

The Alaoglu–Erdős Integer-Exponent Conjecture

Machine-Checked Partial Results, Exact Conditional Structure,
Determinant Methods, and the Remaining Barrier

A unified report on the corpus in

<https://github.com/VladimirReshetnikov/ProveIt>

Bottom line. The conjecture — that $2^x, 3^x \in \mathbb{Z}$ forces $x \in \mathbb{Z}$ — remains open, and the reduction to the four exponentials conjecture explains why no routine combination of present methods can close it. What the corpus does establish is unusually complete. Kernel-verified in Lean: a zero-free region $2^x < 4096$; transcendence of every hypothetical counterexample; multiplicative rank at most two with a canonical primitive odd core; and — conditional on a single counterexample — the *exact* classification

$$\mathcal{S} = \{n + k\beta : n, k \in \mathbb{N}_0\}, \quad \mathcal{C} = \{2^n M^k : n, k \in \mathbb{N}_0\},$$

with unique coordinates, exact dyadic block counts, a sharp quadratic cumulative asymptotic, and vanishing block Fourier modes. These reduce the problem to one narrow statement: no primitive odd $w > 1$ has a positive power of w^ϑ in $\mathbb{Z}[1/3]$.

This report adds, at paper level, two determinant methods — an arbitrary-node Vandermonde repulsion inequality and a generalized-Vandermonde hierarchy over the mixed exponent semigroup $\mathbb{N}_0 + \mathbb{N}_0\vartheta$, the latter giving an independent explicit $O((\log X)^2)$ dyadic bound — together with a diagnosis of exactly why they stop short: archimedean interpolation needs $\asymp (\log X)^2$ nodes at scale X , while a counterexample supplies only $\asymp \log X$. It also records a new sufficient condition of a different kind: if the products $\log p \log q$ of two prime logarithms are $\mathbb{Q}$ -linearly independent, the conjecture follows. The independent continuations reconciled from [Article/new-reports/](#) add a kernel-verified restricted matrix-coefficient equivalence and an exact sparse-interpolation obstruction: R positive points on a hypothetical irrational power curve require exactly $R + 1$ monomials. Its Newton-polygon and boundary-defect refinements isolate a concrete one-monomial saving problem, but do not solve it. The other independent continuations add exact rational-orbit collision counts, a universal superfactorial determinant divisor, a signed-progression denominator-tax identity, and a random-phase/Fejér control experiment. The seventh and eighth independent reports add, respectively, compact/profinite limit laws with exact adelic drift, and an arithmetic audit of polynomial and Padé approximation: denominator clearing, an all-degree normalization-tax obstruction, and rational contact rigidity. Reports 9–13 add de-integralized common divisors and natural boundaries, a fixed integral-matrix and Loewner formulation, finite-place/Fermat-quotient transversality, and a microscopic prime-switch analysis, followed by an exterior-power and boundary-displacement matrix audit. Reports 14–20 add, independently, a structural-prime boundary-resultant package, an information-theoretic audit of the synchronized semigroup, a special-function transfer and mixed-Mahler audit, and a weighted Smith-profile/Loewner hierarchy. They include an exact ordinary/power-sum decoder, algebraic-decoder degree growth, a generic algebraic-base spectrum, a finite mixed-Mahler identity, multilevel modular divisors, and exact geometric-Vandermonde Smith factors, and a Mellin-spectral/approximation audit of the exact conditional monoid. The last two reports add an arithmetic approximation/minimax audit with an exact geometric Newton shadow, and an operator-theoretic classification of which local, finite-lattice, affine, first-jet, and finite-realization extensions would actually add new information. Each sharpens a finite interface while leaving the same four-exponentials barrier intact. Reports 21–27 add independent moment, rank-geometry, unimodular-gap, orbit-entropy, Padé-defect, carry-language, and canonical-product audits. Their formal cores certify exact denominator, matrix, collision, entropy, defect, carry-word, and D-finite identities; their strongest paper endpoints still require a new transcendence, height, uniformity, or interpolation-regularity bridge.

Abstract

Let $\vartheta = \log_2 3$. The Alaoglu–Erdős conjecture asserts that a real x with $2^x, 3^x \in \mathbb{Z}$ must be an integer; equivalently, the only positive integers m with $m^\vartheta \in \mathbb{Z}$ are the powers of two. It is the first unresolved 2×2 instance of the four exponentials problem, while its three-base analogue follows from six exponentials and is formalized here by an interpolation determinant.

This unified report has five purposes. First, it catalogues the machine-checked corpus with exact Lean identifiers: zero-free regions, the Gelfond–Schneider dichotomy, multiplicative rank two, the canonical primitive odd core, the rational-power index and radical degree, localized-radical equivalences, arbitrary-order arithmetic-progression obstructions, shrinking-target Diophantine bounds, fractional-part gap equivalences, and the exact conditional structure theory — the free monoid $\mathcal{S} = \mathbb{N}_0 + \mathbb{N}_0\beta$ with a unique least generator, exact block counts, and the sharp quadratic asymptotic. Second, it develops two new determinant methods. Arbitrary candidate tuples on the lattice curve $y = x^\vartheta$ satisfy

$$\prod_{0 \leq i < j \leq r} (m_j - m_i) \geq \frac{r!}{|(\vartheta)_{\underline{r}}|} m_0^{r-\vartheta},$$

with no arithmetic-progression hypothesis; and the mixed exponent semigroup supplies positive integer generalized Vandermonde determinants whose repulsion exponent tends to one. The paper proof yields the explicit unconditional bound $\#(\mathcal{C} \cap [X, 2X]) \leq 10000 \max\{1, \log X\}^2$ by a route independent of the prime-valuation rank theorem; the corresponding accepted Lean endpoint carries the explicit hypothesis **NewtonFactorization**. On the nonarchimedean side, the accepted corpus now identifies the exact Beatty minimizing fiber and gap, groups every tropical layer by row partitions, truncates the normalized determinant exactly modulo p^K , bounds the number of surviving row partitions by a stars-and-bars coefficient, and extracts finite first-fiber divisors. In the block-scaled consecutive-power model, the all-layer unit gain adds directly to the tropical minimum, giving $p^{\tau+U} \mid \det$ uniformly across the assignment cascade. These gains are lower order; exact synthetic experiments expose genuine cross-layer carrying but do not exhibit a leading-order cascade at the tested sizes. Third, it reconciles the genuinely new material independently developed in all twenty-seven reports under **Article/new-reports/**. The reports share this unified document as their common base; they are audited independently and reconciled by mathematical content rather than treated as a sequential chain. Reports 1–2 show that the Alaoglu–Erdős conjecture is equivalent to a restricted 2×2 Matrix Coefficient principle, while an irrational power curve through R positive nodes has exact sparse interpolation threshold $R + 1$. The corresponding formal prime-symbol determinant, Newton-polygon calibration, and boundary-dilation identities expose an exact factor-four auxiliary-polynomial gap and a finite one-monomial target. Reports 3–5 add the exact rational-orbit collision spectrum, universal factorial and Smith divisors, adelic determinant accounting, and a signed-progression denominator tax. Report 6 adds a uniform metric estimate, exact random-phase extinction, and a positive-kernel Fourier target for the exceptional phase zero. Report 7 adds triangle–Haar and profinite Beatty laws, Sturmian covariance, exact denominator drift and p -adic escape, and a fixed-compact-factor no-go theorem. Report 8 adds exact denominator-clearing identities for rational approximation, a finite-difference normalization-tax obstruction valid in every degree, and the rational contact-equals-multiplicity principle for irrational power curves. Report 9 adds an exact de-integralized common-zero lattice, a Pólya–Carlson boundary, and a compact $\Theta(\log H)$ height dichotomy. Report 10 gives a fixed 2×2 integral-power reformulation and denominator-sensitive Loewner certificates, together with their optimized one-half barrier. Report 11 proves structural-prime p -adic transversality and isolates the Fermat-quotient first layer and adelic error wedge. Report 12 returns to the historical colossally abundant prime-switch problem and shows that simultaneous switches require a microscopic prescribed-prime resonance far below modern short-interval scales. Report 13 develops the matrix direction independently: projective arithmetic density, determinantal-divisor and exterior-power ceilings, and an exact rank- K boundary compression of a K^2 -point spectrum. Report 14 develops the finite-place direction independently: minimal interpolation denominators, structural-prime residual scaling, Newton polygons, and odd-degree dilation resultants in two

favorable residue branches. Report 15 develops an independent information-theoretic direction: common digit-length and Sturmian channels, fixed-rank S -unit sum entropy, exact synchronized-sum decoding, algebraic-decoder degree growth, and a finite Fourier formulation of tied-term cancellation. Report 16 independently tests gamma, Hurwitz, polylogarithmic, Mahler, Barnes, double-sine, and q -digamma transfer mechanisms. It promotes the algebraic-base spectrum to arbitrary independent positive algebraic pairs and isolates a corrected double-pole resonance without confusing resonance with integrality. Report 17 independently replaces the determinant endpoint by a weighted Smith profile at every exterior degree, derives a universal multilevel modular divisor and an exact geometric-Vandermonde Smith form, and couples this arithmetic profile to rational cross-Loewner singular scales. Its entropy analysis proves that a hypothetical semigroup is already extremal for ordinary anti-concentration and Shannon inequalities. Report 18 independently reads the exact monoid as a diagonal spectrum: rank one and rank two become simple and double spectral-zeta poles, linear and quadratic logarithmic Weyl laws, and one or two Mellin frequency lattices. It also proves a sharp fixed-window polynomial node-budget barrier and isolates an exact two-scale support equation, while showing that none of these formal rank signatures uses the decisive extra integrality $M^\vartheta \in \mathbb{N}$. Report 19 independently develops the corresponding arithmetic approximation lattice: an exact unit-locking principle, the same optimized fixed-window capacity, finite dyadic interpolation certificates, a cofactor formula for discrete minimax error, and a geometric Newton expansion whose first omitted value is explicit. Report 20 independently audits the operator interface: two-mode shifts, heat data, and semisimple functional calculus are blind, whereas a global or finite valuation lattice, one affine coherence, one rational first jet, or one exact finite trace/Hankel realization would force integrality. Report 21 adds a Hausdorff moment representation, exact structural-prime moment defects, finite rational-quadrature failure, and Müntz-invisible positive perturbations. Report 22 classifies the depth-two logarithmic tensor by hyperbolic rank and supplies an exact rational symmetrization. Report 23 proves determinant-distance gap gcds, transported rational secants, a finite order-signature no-go, and an entrywise jet ceiling. Report 24 computes the conditional orbit heights exactly and turns Farey counting into simultaneous worst-case denominator entropy. Report 25 optimizes Kawashima’s projective Padé defect and proves that no positive-defect rational tuple can recover the fixed prime logarithm. Report 26 develops distinguished-descendant carry languages, exact multinomial fibers, and their finite and unrestricted pressure laws. Report 27 identifies the sharp cubic logarithmic type of the 3-smooth canonical product, proves subcritical interpolation rigidity, and isolates an exact critical boundary cocycle. Fourth, it quantifies the remaining gap: the determinant hierarchy misses the conditional population by exactly one factor of $\log X$. Fifth, it records new equivalent and sufficient conditions — block-growth dichotomies (including the observation that the conjecture is equivalent to the dyadic block count being 1 for *infinitely many* exponents), a compact synchronized S -integral orbit formulation, and the prime-log-product independence criterion — and closes with a prioritized research program and a formalization blueprint.

Contents

1	Executive summary	1
1.1	What is proved, and at what level of trust	1
1.2	Headline results	2
1.3	The strongest new conclusions	4
1.4	Why this is not a proof	5

2	The problem and its equivalent forms	5
2.1	Original statement	5
2.2	Candidate-base normalization	6
2.3	The logarithmic determinant	7
2.4	One-parameter subgroup form	7
2.5	The four-exponentials reduction	7
2.6	Rational and algebraic exponents	8
2.7	The three-base theorem and its determinant architecture	9
3	Repository timeline and verification status	10
3.1	Evidence convention	11
4	The kernel-verified corpus	11
4.1	Fractional-part gaps and shrinking targets	11
4.2	Verified zero-free regions and certificate architecture	12
4.3	Arithmetic descent	13
4.4	Local progression obstructions and zero density	14
4.5	Multiplicative rank two and the common odd core	15
4.6	Transcendence	16
4.7	Classification of auxiliary bases: the 3-smooth theorem	16
4.8	Unconditional gap equivalences	16
5	Exact conditional structure of all solutions	17
5.1	The primitive common odd core	18
5.2	The rational-power index and the denominator normalization	19
5.3	The exact primitive-generator theorem	19
5.4	Primitivity of the least output pair	21
5.5	Exact counting	22
5.6	Blockwise irrational rotation and equidistribution	23
6	Equivalent one-number and radical formulations	24
6.1	Localization at the prime three	24
6.2	Reduction to power-primitive odd bases	25
6.3	Canonical radical algebraicity	25
6.4	Where the wall stands in the smallest case	26

7	The exact four-exponentials barrier	27
7.1	The two-by-two exponential matrix	27
7.2	Why six exponentials does not specialize	27
7.3	Integrality is stronger than algebraicity, but not by enough	28
8	Rational-function rigidity and closure properties	28
8.1	Evaluation injectivity at the generator	28
8.2	The rational-function intersection theorem	29
8.3	Exact multiplication criterion	30
8.4	Exact quotient criterion	31
8.5	Closure formulations of the conjecture	32
8.6	Formalization cost	33
8.7	What the rigidity statements buy	34
9	Three-smooth outputs and transcendental product exponents	34
9.1	The two outputs cannot both be three-smooth	35
9.2	Third-base transcendence from six exponentials	36
9.3	Products of noninteger solutions	37
9.4	A global second-level trichotomy	38
9.5	Algebraicity formulations equivalent to the conjecture	39
9.6	Formalization boundary	40
10	Perfect-power content, affine descent, and primitive-output statistics	41
10.1	Simultaneous perfect powers	41
10.2	Matched affine decompositions	42
10.3	What is already formal, and what is new here	43
10.4	Why this does not yet descend below β	44
10.5	Total and primitive counts	45
10.6	Distribution of the maximal common power degree	46
10.7	Exact generating series	46
10.8	Interpretation	47
11	Candidate divisibility and ambient root saturation	47
11.1	A two-dimensional valuation cone	47
11.2	The candidate gcd/lcm lattice	48

11.3	Fixed prime support	49
11.4	Ambient root saturation and the index d	50
11.5	Why lattice saturation is a good Lean target	52
12	A sufficient condition: independence of prime-log products	53
12.1	Failure as a vanishing linear form in log-products	53
12.2	The independence criterion	53
12.3	What the criterion is, and is not	54
12.4	A finite form of the criterion	56
12.5	The two-prime case is a theorem, and it is sharp	56
13	The algebraic auxiliary-base spectrum	57
13.1	The saturated group of positive two-three units	57
13.2	The exact spectrum theorem	58
13.3	The divisible-hull formulation	59
13.4	The determinant engines behind the spectrum	60
13.5	The strong spectrum via Baker and Roy	61
13.6	One base is enough: detectors and rank equivalences	61
13.7	Exact exponent, rational-function, and tower loci	62
13.8	Why the exact spectrum does not settle the conjecture	65
14	Rational and integral outputs inside the spectrum	66
14.1	The valuation matrix of an output pair	66
14.2	Smith normal form and the finite rational-output defect	67
14.3	Two translated solutions remove the fractional defect	70
14.4	The two smooth branches are mutually exclusive	71
15	The second-iterate kernel and multiplicative dependence	72
15.1	The kernel and its invariance	72
15.2	Multiplicative dependence of four integers	73
15.3	Why the multiplicative form is worth having	74
16	Depth-three tower rigidity	75
16.1	The saturated spectrum and its strong form	75
16.2	The return kernel	77

16.3	Exact survival depth	78
16.4	Exact behaviour of the bases two, three, and six	79
16.5	One-base tower equivalences	80
17	Symmetric tower constants	82
17.1	An unconditional strong dichotomy	82
17.2	Exact conditional classification	83
18	Finite-difference obstructions	85
18.1	The sharp three-term criterion	86
18.2	Thresholds for higher order	86
18.3	Why longer progressions do not close the problem	86
18.4	The optimized variable-order scale	87
19	New result I: arbitrary-node lattice repulsion	90
19.1	The determinant identity	90
19.2	The arbitrary-node theorem	90
19.3	Three points: a lattice-triangle inequality	91
19.4	General local packing bound	92
19.5	A bound for any longer interval	93
19.6	Relation to higher differences	93
20	New result II: semigroup generalized Vandermonde determinants	94
20.1	Positivity of generalized Vandermonde determinants	94
20.2	Integer determinants from candidates	94
20.3	Divided-difference factorization	95
20.4	Quantitative semigroup gap theorem	96
20.5	The ordered exponent semigroup	96
20.6	An explicit coarse form	97
20.7	Counting candidates in a general interval	98
20.8	An independent dyadic logarithmic-square bound	98
21	The min-plus ceiling: why term-wise p-adic amplification cannot close a determinant argument	99
21.1	The threshold	99
21.2	The ceiling	99

21.3	A correction to the averaging form of the obstruction	100
21.4	Calibration: the exact ceilings	101
21.5	Where cancellation can live	102
21.6	What this leaves	107
22	Why the new hierarchy still stops short	108
22.1	How many nodes are needed to approach scale X ?	108
22.2	How many nodes a counterexample supplies	108
22.3	Broader multiplicative intervals do not help	109
22.4	The determinant needs one more source of size	109
23	Further equivalent reformulations	110
23.1	Integer points on a transcendental lattice curve	110
23.2	Vanishing determinant of four logarithms	110
23.3	Primitive localized-radical exclusion	110
23.4	Exact growth dichotomy	110
23.5	Compact synchronized S -integral orbit	112
23.6	Two-dimensional exponential grid	112
23.7	An operator formulation on the prime-logarithm space	113
24	A consolidated catalogue of equivalent formulations	114
24.1	How to read the catalogue	114
24.2	Cluster A: normalized restatements	114
24.3	Cluster B: localized radicals	115
24.4	Cluster C: counting and block growth	116
24.5	Cluster D: orbits and fractional parts	117
24.6	Cluster E: closure, smoothness, and algebraicity	118
24.7	Which formulations are worth attacking	118
24.8	Equivalences versus sufficient conditions	119
25	Independent finite verification	120
25.1	What each scan must certify	120
25.2	Superseded scans	121
25.3	The trust boundary, stated plainly	121

26 Attempts at a full proof and their exact failure points	122
26.1 Collapsing rank two to rank one	122
26.2 Rational approximation to a least generator	122
26.3 Equidistribution against exponentially shrinking targets	122
26.4 Finite differences and additive progressions	123
26.5 Radical fields and conjugates	123
26.6 The compact S -integral orbit	124
26.7 Forcing an auxiliary base	124
26.8 Exploiting the primitive output pair	124
26.9 The linear-form-in-log-products formulation of the wall	125
27 Consolidated research program	125
27.1 Priority A: p -adic amplification of semigroup determinants	125
27.2 Priority B: exploit the full (n, k) grid, not sorted magnitudes	126
27.3 Priority C: denominator-sensitive counting on the compact arc	126
27.4 Priority D: mixed archimedean–radical invariants	126
27.5 Priority E: discharge and sharpen the determinant hypotheses	127
27.6 Priority F: force a third algebraic exponential	127
27.7 Priority G: five and sharp six exponentials	127
27.8 Priority H: larger kernel-replayed regions and effective measures	127
27.9 Priority I: computational reconnaissance with exact certificates	128
27.10 Priority J: the colossally abundant interface	128
28 Lean formalization blueprint	128
28.1 ArbitraryNodeDeterminant	129
28.2 GeneralizedVandermonde	130
28.3 ExponentialPolynomialZeros	130
28.4 MixedExponentSemigroup and SemigroupDeterminant	131
28.5 FallingFactorialBound	132
28.6 SemigroupWeightBound	132
28.7 SemigroupDyadicBound	133
28.8 BlockGrowthReformulations	133
28.9 CompactOrbit	134
28.10 LogProductIndependence	134

28.11	Proof-audit discipline	135
29	Comparison of methods	135
29.1	Prime-valuation rank versus interpolation geometry	135
29.2	Exact conditional rank one over two generators	136
29.3	Finite differences as confluent determinants	136
30	Conclusions	137
31	Continuation: executive answer, scope, and trust boundaries	139
31.1	The requested goal and the outcome	139
31.2	The principal new results	139
31.3	Evidence labels used in this part	140
31.4	What the determinant ceilings do and do not prove	140
32	Baseline inherited from the unified report	141
32.1	Notation	141
32.2	Inherited results used here	141
32.3	Repository reconciliation as of August 21, 2026	142
32.4	The exact remaining logarithmic wall	142
33	A synchronized Sturmian model of a counterexample	143
33.1	Fractional-part normalization	143
33.2	Sturmian consequences	143
33.3	Exact reduced denominators	144
33.4	An exact equivalent exclusion principle	144
33.5	Cobham-type and Mahler-type diagnostics	145
34	Effective logarithmic separation	146
34.1	From linear forms in two logarithms to an irrationality measure	146
34.2	Polynomial recurrence for convergent denominators	146
34.3	Polynomial spacing in conditional blocks	147
34.4	Separation on the full (n, k) lattice	147
34.5	Why the Baker upgrade does not close the conjecture	148
35	Exact all-support valuations of ordinary block Vandermondes	148

35.1	The block determinants	148
35.2	The input determinant: an exact formula at every support prime	149
35.3	The output determinant and bounded tie corrections	150
35.4	The depth-sensitive product ceiling	151
35.5	Why even 87% is not enough	151
36	The mixed-semigroup determinant and its Gini norm	152
36.1	Balanced low-weight columns	152
36.2	The all-prime tropical divisor	153
36.3	Antimonotone assignment and the Gini norm	153
36.4	Closed form of the Gini norm	154
36.5	Normalized prime-valuation slope vectors	154
36.6	The exact continuum limit	155
36.7	The universal $4/5$ ceiling	156
36.8	Cross-coprime constants: $7/10$ and $3/10$	157
36.9	Correction and calibration of a later repository note	158
36.10	Where extra divisibility can occur	158
37	The complete triangular grid and a $\pi/4$ transport ceiling	159
37.1	Rows and columns	159
37.2	The common triangular distribution	160
37.3	Summing all primewise assignment minima	160
37.4	Optimal transport and the $\pi/4$ ratio	161
37.5	The Harish–Chandra–Itzykson–Zuber identity	162
38	The first tied p-adic layer is lower order	163
38.1	Tropical initial form at two	163
38.2	The symmetric initial form at three	164
38.3	Two-logarithm bounds inside the fibers	165
38.4	What a leading excess would have to do	166
38.5	Exact finite cascade reconnaissance	166
39	Synthesis of the determinant audits	167
39.1	Correction of the post-attachment “one half / two thirds” note	168

40 Directed-rounding extension through 2^{27}	169
40.1 Manifest audit	169
40.2 The new range	170
40.3 New extremal records and high-precision replay	170
41 Additional proof routes pushed to their failure points	171
41.1 Transferring convergents of β to ϑ	171
41.2 Two-row determinants and fixed-support S -units	172
41.3 Automaticity and the synchronized Sturmian word	173
41.4 Real and p -adic logarithms	173
41.5 Transcendental-curve point counting	174
42 Exact sufficient conditions extracted from the new analysis	174
42.1 The mixed-block cancellation-excess criterion	174
42.2 A denominator-cancellation criterion	175
42.3 An automaticity criterion	175
43 Branch-sensitive research program	176
43.1 Branch I: independent external valuation directions	176
43.2 Branch II: a common external valuation ray	176
43.3 Branch III: one smooth output	177
43.4 Branch IV: cross-coprime primitive outputs	177
43.5 Archimedean priority: solve the triangular HCIZ limit	177
43.6 Finite verification priority	178
44 Lean formalization blueprint	178
44.1 Proposed module graph	178
44.2 Finite statements before asymptotics	180
44.3 Exact theorem sketches	180
44.4 The formalized finite core of the first tied layer	181
44.5 Certificate architecture for the new finite bound	181
45 Numerical sanity checks and reproducibility	182
45.1 Balanced-block assignment constants	182
45.2 Complete-grid transport ratio	182

45.3 Directed-rounding scan manifest	183
45.4 Extremal replay records	183
46 Status matrix and dependency audit	184
46.1 Logical dependency graph	186
47 Conclusions	186
48 Second continuation: scope and claim ledger	188
48.1 Claim ledger	188
49 Effective Diophantine and counting consequences	189
50 Valuation-lattice saturation and an exact duality	190
51 Exact-grid interpolation determinants	191
51.1 Nonvanishing and total positivity	192
51.2 Archimedean upper bound	192
51.3 Finite-prime lower bound as an assignment problem	192
52 Square grids: a universal divisor and the critical cancellation	193
53 Triangular grids and a new exclusion inequality	194
53.1 Limiting distributions	194
53.2 The archimedean constant	194
53.3 The finite-prime constant	195
53.4 Numerical anatomy of the two grid geometries	196
54 Rational changes of basis and the product formula	196
54.1 A signed arithmetic functional	196
54.2 All unimodular bases of $\langle 2, 3 \rangle$	197
54.3 The basis $(2, 3/2)$ and a gcd constraint	198
55 Joint prime transport and shared-factor gains	199
56 Universal congruence divisibility is lower order	200
57 Why several tempting routes still stop short	200

57.1 Continued fractions and integer gaps	200
57.2 Catalan, <i>abc</i> , and small differences	200
57.3 Counting rational points	201
57.4 The four-exponentials boundary	201
58 A variational reformulation of the remaining determinant problem	201
59 Confluent and integer-valued determinant extensions	202
60 Conditional resolutions and exact reformulations	202
60.1 Four Exponentials and Schanuel	202
60.2 Logarithm determinant formulation	202
60.3 Order-preserving monoid formulation	203
61 Formalization roadmap	203
62 Prioritized research directions	203
63 A proof-producing finite verification	204
64 Conclusion	205
65 Third continuation: matrix coefficients, sparse interpolation, and boundary defects	205
65.1 The restricted matrix-coefficient principle	206
65.2 The formal prime-symbol determinant	207
65.3 Exact sparse interpolation on an irrational power curve	207
65.4 BK degree and the exact factor-four threshold	208
65.5 Boundary defects of the extremal interpolant	209
65.6 Counting thresholds and rational monodromy	209
65.7 The reported 2^{29} scan and its trust boundary	210
65.8 What changed, and what did not	210
66 Further continuations: support spectra and arithmetic determinant audits	211
66.1 The exact orbit-collision support spectrum	211
66.2 Coefficient height and the generic-modulus no-go	212
66.3 Universal factorial content and primitive interpolation	212

66.4	The adelic–HCIZ accounting identity	213
66.5	Signed progressions and the denominator tax	214
67	Sixth report: phase randomization and positive-kernel certificates	215
67.1	A uniform metric estimate and random-phase extinction	216
67.2	The Fejér certificate	216
67.3	Why naive random sparse coefficients do not bridge the gap	217
68	Seventh and eighth reports: compact laws and approximation rigidity	218
68.1	Difference groups, triangle–Haar laws, and profinite Beatty factors	218
68.2	Exact denominator drift and noncompact escape	219
68.3	Sturmian rigidity and fixed-observable blindness	219
68.4	Arithmetic resolution of normalized polynomial approximation	220
68.5	Rational contact rigidity and the area-to-diameter collapse	221
69	Ninth report: analytic common divisors and natural boundaries	222
69.1	Removing the automatic zeros and counting the conditional divisor	223
69.2	The synchronized word has a natural boundary	223
69.3	A logarithmic compact-height criterion	224
70	Tenth report: matrix powers, Loewner certificates, and the one-half barrier	225
70.1	One fixed integral matrix	225
70.2	Cross-Loewner and generalized-alternant certificates	226
70.3	Quantitative repulsion and the optimized barrier	226
71	Eleventh report: finite-place transversality and local rank-drop certificates	227
71.1	The structural primes are locally transverse	228
71.2	The Fermat-quotient first layer	228
71.3	Real approximants and the exact adelic wedge	229
71.4	Cascade depth and the Kummer norm diagnostic	229
72	Twelfth report: prime switches and microscopic resonance	230
72.1	Exact switch heights and logarithmic means	230
72.2	A microscopic prescribed-prime window	231
72.3	The exact 2–3 convergent sieve	231
72.4	What prime distribution can and cannot force	232

73 Thirteenth report: exterior powers and boundary realizations	233
73.1 Rational conditioning does not create rational rank	233
73.2 Projective arithmetic density and a Hadamard-power no-go	234
73.3 Determinantal divisors and the compound ceiling	234
73.4 Exact boundary displacement and transfer determinants	235
73.5 Kraus–Loewner positivity and its denominator tax	236
74 Fourteenth report: structural-prime residuals and dilation resultants	237
74.1 Local rank is thick, but it is not a product-formula input	237
74.2 The interpolation denominator is a cokernel order	238
74.3 Exact scaling, roots, and an honest algebraic resultant	238
74.4 Paired residuals and the full structural Smith profile	239
75 Fifteenth report: synchronized information and decoder complexity	245
75.1 One magnitude channel and a rank-two valuation profile	245
75.2 Fixed-rank sums are asymptotically unordered-pair codes	245
75.3 Exact synchronized decoding and degree growth	246
75.4 Fourier separation identifies, but does not solve, tied cancellation	246
76 Sixteenth report: transfer spectra and mixed Mahler resonance	247
76.1 The algebraic-base spectrum	247
76.2 Entire cancellation and rank-preserving transfer formulas	248
76.3 A finite Mahler identity and the corrected resonance theorem	248
76.4 Double-sine, q -digamma, and the surviving target	249
77 Seventeenth report: weighted Smith profiles and rational Loewner clusters	250
77.1 From one determinant to every exterior level	250
77.2 Multilevel residue collapse and a geometric Smith model	250
77.3 Rational cross-Loewner clusters	251
77.4 Positivity and entropy no-go theorems	252
78 Eighteenth report: logarithmic spectral rank and approximation capacity	253
78.1 Spectral zeta and logarithmic Weyl rank	253
78.2 Complete heat-trace expansion	253
78.3 The finite two-scale support equation	254

78.4	A sharp fixed-window node-budget barrier	255
79	Nineteenth report: arithmetic approximation and geometric Newton interpolation	255
79.1	Subunit locking and the optimal multiplicative block	256
79.2	Dyadic certificates and the discrete minimax unit barrier	256
79.3	The geometric Newton shadow	257
80	Twentieth report: operator coherence and local-to-global arithmetic	258
80.1	Universal two-mode dynamics	258
80.2	Global and finite arithmetic covariance	258
80.3	Affine rulers and the first-jet criterion	259
80.4	Sturmian similarity and exact finite realization	259
81	Varying the antecedent, the domain, and the codomain	260
81.1	The antecedent ladder	260
81.2	Complexifying the exponent is free	261
81.3	The algebraic antecedent over the complex numbers	262
81.4	Multivalued branches widen the wall by one prime	262
81.5	Complexifying the codomain: the Gaussian version is equivalent	263
81.6	Summary	263
82	Twenty-first report: Hausdorff moments and Müntz invisibility	264
82.1	The explicit moment measure	264
82.2	Exact arithmetic defects and finite rational simulation	264
82.3	All smooth linear observations remain invisible	265
83	Twenty-second report: depth-two rank geometry and rational symmetrization	266
83.1	Baker excludes quadratic rank at most two	266
83.2	Hyperbolic pullback and the rank-three conic	266
83.3	A rational symmetric normal form	267
84	Twenty-third report: unimodular gaps, transported secants, and finite-order barriers	267
84.1	Determinant distance and adjacent gap coprimality	267
84.2	Rational secants and separated denominators	268

84.3	Finite signatures and the entrywise jet ceiling	268
85	Twenty-fourth report: exact orbit heights and Farey entropy	269
85.1	Exact reduced orbit law	269
85.2	Worst-case Farey-to-orbit distortion	270
85.3	Baker recurrence versus fixed contact	270
86	Twenty-fifth report: projective Padé defect and rational-transfer no-go	271
86.1	The projective defect bound	271
86.2	No rational transfer of the prime direction	272
87	Twenty-sixth report: carry languages and fixed-target entropy	273
87.1	Carry paths and finite alphabets	273
87.2	Exact carry-language pressure	274
88	Twenty-seventh report: canonical-product threshold and critical defect cocycle	275
88.1	Regularity ladder for a power branch	275
88.2	The sharp cubic logarithmic type	276
88.3	The critical boundary cocycle	276
A	Status matrix for major statements	277
B	Formalization inventory	289
B.1	Kernel-accepted modules	289
B.2	Former theorem drafts and the remaining finite-certificate draft	295
C	A symmetric formulation	296
D	Detailed proof of the generalized Vandermonde sign	297
E	Numerical constants	298
F	Interval computation: script and outputs	298
G	Suggested experimental determinant families	300
G.1	Rectangular row and column boxes	300
G.2	Triangular boxes adapted to height	301
G.3	Finite-difference row operations	301

G.4	Confluent columns	301
H	Dependency map and Lean-oriented proof sketches	301
H.1	Dependency map	301
H.2	Multiplication skeleton	302
H.3	Output smoothness skeleton	303
H.4	Common-power skeleton	304
H.5	Ambient saturation skeleton	304
H.6	Suggested theorem order for a pull request	305
I	Directed-rounding verifier	305
I.1	What the verifier certifies	305
I.2	Source identity and environment	306
I.3	Source	306
I.4	Build and invocation	310
J	Chunk audit and reproducibility	310
J.1	Extremal records	311
J.2	Reproducibility checklist	312
K	Classical transcendence inputs and their exact roles	312
K.1	Logical dependency graph	315
L	Common false shortcuts and audit warnings	316
M	Logarithmic-energy convergence for projected lattice triangles	317
N	Integral evaluations used in the determinant constants	318
N.1	The beta-two logarithmic energy	318
N.2	The simplex assignment constant	318
N.3	The balanced-rectangle energy	319
O	General two-prime version	319
P	Finite determinant certificates	319
Q	Source of the finite verifier	320

R	External literature reconnaissance	322
R.1	Effective measures for the two logarithms	322
R.2	Interpolation determinants with non-archimedean amplification	323
R.3	Small-value estimates and the four-exponentials frontier	323
R.4	Counting rational points on the transcendental arc	324
R.5	Multiplicative dependence, S -units, and rigidity	324
R.6	Kummer theory and the radical reformulation	324
R.7	Why the generic-power model theory excludes this exponent	325
R.8	Prioritized directions, and what the ceiling changes	325
R.9	Provenance of the conjecture	326
R.10	The formalization landscape	326
R.11	Summary of what the literature does not supply	326

1 Executive summary

1.1 What is proved, and at what level of trust

The full conjecture remains open. Several routes were pushed to their exact failure points: the three-base determinant method does not shrink to two bases without proving a case of the four exponentials conjecture; equidistribution of multiples does not defeat exponentially shrinking integrality gaps; local curvature obstructions cannot reach the sparse global structure; rational approximation to a hypothetical least generator stalls at the same transcendence input; and every successful reduction terminates at the same transcendence-theoretic wall, isolated in Section 6. Two lines of attack opened after that survey – interpolation determinants on arbitrary candidate nodes, and generalized Vandermonde determinants over the full semigroup of integral mixed monomials – do produce new unconditional partial results, and a third, the prime-log-product criterion of Section 12, isolates a clean sufficient condition of a different flavor. None of them closes the final gap.

To keep unlike kinds of evidence separate, every important statement in this report carries one of five labels.

[kernel-verified Lean]

A proof source is present in the repository, lies on the audited `ExponentialIdentities.lean` import surface, and has been accepted by the Lean kernel in direct or full builds of the cited modules (toolchain 4.31/4.32 series, Mathlib pinned by the repository manifest). No `sorry`, no extra axioms, no `native_decide`.

[paper proof in this report]

A conventional mathematical proof is supplied in this document. It is intended to be suitable for formalization, but no claim of kernel checking is made.

[Lean draft, not yet accepted]

A Lean source for the statement exists in the repository tree, but it has *not* yet been accepted by the kernel at the pinned state. A draft is weaker evidence than [paper proof in this report] prose backed by a written proof, not stronger: it records intent to formalize, nothing more.

[software computation]

A finite statement was checked by an explicitly identified software computation, usually with outward rounding. The local section records the script, parameters, margins, artifacts, and relevant trust base: this includes CPython and `mpmath` for the interval checks and C, MPFR, and GMP for the directed-rounding scans. Such evidence is not interchangeable with a Lean certificate.

[classical/open background]

A classical theorem quoted from the literature.

An earlier audit, performed against 423ac24, correctly observed that three modules cited by the survey — `FiniteCheck4096`, `OddCore` and `FractionalPartDensity` — were absent then. That

discrepancy is resolved: the three modules were subsequently committed in 98646b455 and are kernel-verified; the survey’s temporary pending-verification remark was removed in 716617d72. Statements from these modules are therefore labeled [\[kernel-verified Lean\]](#) below. The kernel-checked feature wave also includes

PrimitiveOddCore, LeastSolution, ThreeDenominatorNormalization,
 RationalPowerIndex, PrimitiveGenerator, ExactCounting,
 QuadraticAsymptotic, BlockWeyl, BlockEquidistribution,
 OutputNormalization, RationalThirdBase, RadicalReformulation,
 PrimitiveRadical, RadicalDegree, SharperProgressions,
 ThirdDifference, HigherDifference.

The determinant and reformulation modules that earlier editions of this report listed as drafts have since been accepted, and a further round added

CandidateLattice, KernelDichotomy, SecondIterateKernel,
 TowerRigidity, SymmetricTowerConstants, PrimeLogSpaceOperator,
 ArbitraryNodeRepulsion, SemigroupGapBound.

The current algebraic-locus and nonarchimedean wave further includes

AlgebraicExponentLocus, AlgebraicOutputLocus,
 RationalFunctionAlgebraicLocus, AlgebraicTowerTransition,
 CrossBaseTowerTransition, BlockVandermondeValuationBudget, AllLayerUnitBasis,
 UnitOrderingGain,
 SimultaneousRadicalDegrees, CanonicalRadicalMixedOutput,
 TropicalInitialForm, BlockFiberFactorization, TropicalBlockApplication,
 BeattyTropicalFiber,
 RowBlockLaplaceExpansion, AssignmentExcessGeometry, RowBlockUnitDivisibility,
 MixedDeterminantUnitBridge,
 RestrictedMatrixCoefficient, FormalPrimeDeterminant, SparsePowerCurve,
 RationalOrbitGrid, BoundaryDilationDefect, SuperfactorialDeterminant,
 SignedProgressionDeterminant, FejerKernelCertificate, GlobalMixedPrefixCascade,
 DilationDifferenceQuotient, SynchronizedSumDecoder, RationalSumDecoder,
 AlgebraicBaseSpectrum, MixedMahlerFinite, GeometricVandermondeSmith,
 TwoScaleSupportFinite, GeometricNewtonInterpolation, JordanFirstJet.

Public headline theorems throughout were audited with `##print axioms`; only the standard Mathlib axioms `propext`, `Classical.choice`, and `Quot.sound` occur, and the audit for the most recent round is itself a module, `AuditNewModules`. The two analytic inputs still missing from the determinant chain are carried as named hypotheses of the theorems that use them, not as axioms. The per-statement ledger is [Appendix A](#) and the module inventory is [Appendix B](#).

1.2 Headline results

1. **Verified zero-free region.** If $2^x, 3^x \in \mathbb{Z}$ and $2^x < 4096$, then $x \in \mathbb{Z}$; equivalently every nonintegral solution has $x \geq 12$. [\[kernel-verified Lean\]](#) A directed-rounding scan extends the finite barrier to $2^x < 2^{27}$, i.e. $x > 27$. [\[software computation\]](#)

2. **Transcendence.** Every nonintegral solution is transcendental; for x with $2^x \in \mathbb{Z}$, algebraicity of x is *equivalent* to integrality. $\vartheta = \log_2 3$ is transcendental. [\[kernel-verified Lean\]](#)
3. **Exact conditional structure.** Conditional on a counterexample there is a unique gcd-normalized odd core $w > 1$, every candidate has unique coordinates $2^i w^j$, and the solution set has the exact free-monoid classification $\mathcal{S} = \mathbb{N}_0 + \mathbb{N}_0 \beta$ with a unique least nonintegral generator β , which is affinely primitive. Candidate counts below N are at most $\log_2 N \cdot \log_3 N$, each dyadic block $[2^T, 2^{T+1})$ holds at most $\log_3(2^{T+1})$ candidates, the exact number of *nonintegral solutions* in the exponent block $[N, N+1)$ is $\lfloor (N+1)/\beta \rfloor$, while the full solution count — equivalently the candidate count in $[2^N, 2^{N+1})$ — is $1 + \lfloor (N+1)/\beta \rfloor$, the cumulative count satisfies $C(T)/T^2 \rightarrow 1/(2\beta)$, and the conjecture is equivalent to subquadratic growth along 2^T . Exact block averages of every finite trigonometric polynomial converge to its constant Fourier coefficient. [\[kernel-verified Lean\]](#) The extension to all continuous tests and interval indicators remains paper-level. [\[paper proof in this report\]](#)
4. **Unconditional gap equivalences.** The conjecture is equivalent to the existence of any nonempty open subinterval of $(0, 1)$ free of fractional parts $\{x\}$ of nonintegral solutions, and to those fractional parts being bounded away from 0, or from 1; the convergent denominators of ϑ obey $q_{n+1} < 2m^{q_n}$. [\[kernel-verified Lean\]](#)
5. **Sharper local combinatorics.** For every $r \geq 2$, the sharp criterion $|(\vartheta)_r| h^r < n^{r-\vartheta}$ forbids $r+1$ candidates in arithmetic progression, uniformly in the progression length. The underlying arbitrary-order forward-difference mean-value identity is formalized as well, and for the three-term case the rational effective criterion $h^{400} \leq n^{83}$ (step exponent $83/400 = 0.2075$) is verified independently. [\[kernel-verified Lean\]](#)
6. **New: arbitrary-node repulsion and the semigroup hierarchy.** Any $r+1$ distinct natural candidates — with no arithmetic-progression hypothesis — obey a quantitative product-of-gaps lower bound, giving explicit local packing bounds and a best ordinary-power exponent $0.241503\dots$ at determinant order 4 (Section 19). Enlarging the exponent set to the mixed semigroup $\mathbb{N}_0 + \mathbb{N}_0 \vartheta$ yields arbitrarily large positive integer generalized Vandermonde determinants, a near-linear repulsion exponent $1 - O(r^{-1/2})$, and the paper-level unconditional dyadic bound $\#(\mathcal{C} \cap [X, 2X]) \leq 10000 \max\{1, \log X\}^2$ (Section 20). This route is independent of the repository’s prime-valuation rank theorem. The paper theorems are unconditional. Their current Lean forms isolate the two missing analytic ingredients as the explicit hypotheses `RpowDividedDifferenceMeanValue` and `NewtonFactorization`; all determinant identities, integrality, nonvanishing, positivity, and the deductions from those hypotheses are kernel-accepted. [\[paper proof in this report\]](#) [\[kernel-verified Lean\]](#) Lean modules `ArbitraryNodeDeterminant`, `MixedExponentSemigroup`, `SemigroupDeterminant`, `FallingFactorialBound`, `SemigroupWeightBound`, `ArbitraryNodeRepulsion`, and `SemigroupGapBound` are on the aggregate import surface.
7. **New: a sufficient condition from products of prime logarithms.** Suppose $2^x = M \in \mathbb{Z}$ and $3^x = A \in \mathbb{Z}$. Eliminating x gives the exact identity

$$\log M \log 3 = \log A \log 2.$$

Expanding A and M into prime factorizations turns this into an integer-coefficient linear form in the pairwise products $\log p \log q$ of prime logarithms, and the form is nontrivial precisely when $x \notin \mathbb{Z}$. Hence: *if the products $\log p \log q$, over unordered pairs of primes, are linearly independent over $\mathbb{Q}$, the conjecture holds* (Section 12). That independence is a consequence

of Schanuel’s conjecture. It is different in kind from the four exponentials conjecture – it constrains a bilinear lattice of logarithm products rather than a 2×2 logarithm determinant – so it supplies a genuinely separate route to the same conclusion, not a reformulation of the same wall. [\[paper proof in this report\]](#) [\[kernel-verified Lean\]](#) The exact conditional implication is accepted in `LogProductIndependence`; the asserted prime-log-product independence itself remains a classical hypothesis, not a theorem of the corpus.

8. **Exact localization of the p -adic cascade.** For the mixed Beatty determinant, the minimum assignments, their floor-sized gap, and the first block-determinant quotient are now exact. More generally, after removing p^τ , reduction modulo p^K can receive contributions only from row partitions of assignment excess below K , with every coefficient a signed product of square unit minors. The first-fiber valuation has an elementary finite cube budget, and monic unit-adapted bases extract additional whole-Vandermonde divisors. For the block-scaled consecutive-power model, their common exponent U is independent of the row partition and the exact additive divisor $p^{\tau+U} \mid \det$ follows. These results also bound the number of terms below any fixed excess cutoff by an explicit binomial coefficient. They localize every possible leading excess to the valuations of, and cancellation among, the surviving low-excess minor products; they bound neither phenomenon. [\[kernel-verified Lean\]](#)
9. **New: restricted matrix coefficients and exact sparse interpolation.** For positive integral outputs M, A , the singular logarithm matrix $\begin{bmatrix} \log 2 & \log M \\ \log 3 & \log A \end{bmatrix}$ has a zero coefficient between nonzero rational vectors exactly as required by the restricted Matrix Coefficient principle if and only if the Alaoglu–Erdős conjecture holds. Moreover, on an irrational curve $Y = X^\beta$, vanishing at R distinct positive nodes requires at least $R + 1$ monomials, and a positive-node interpolant attains this bound with every coefficient nonzero. The lower bound, integral evaluation determinant, matching interpolant, and matrix-coefficient equivalence are kernel verified. The induced BK-degree threshold $R - 1$ and boundary compression leave a precise one-monomial saving target. [\[kernel-verified Lean\]](#) [\[paper proof in this report\]](#)
10. **The irreducible remainder.** The conjecture is equivalent to: no odd integer $u > 1$ satisfies $u^\vartheta \in \mathbb{Z}[1/3]$, with a symmetric base-swapped formulation and an equivalent reduction to canonical power-primitive odd bases. The least radical index gives an irreducible binomial and the exact algebraic degree. [\[kernel-verified Lean\]](#) The exclusion itself remains open even for $u = 3$.

1.3 The strongest new conclusions

Determinant hierarchy (Sections 19 and 20). For every $r \geq 2$, any $r + 1$ distinct natural candidates obey a quantitative product-of-gaps lower bound, with no arithmetic-progression hypothesis. Enlarging the exponent set from ordinary powers to the mixed semigroup $\mathbb{N}_0 + \mathbb{N}_0\vartheta$ supplies arbitrarily large positive integer generalized Vandermonde determinants; their archimedean upper bounds give a near-linear repulsion exponent $1 - O(r^{-1/2})$ and, unconditionally,

$$\#(\mathcal{C} \cap [X, 2X]) \ll (\log X)^2,$$

with the explicit constant 10000 proved below. Paper proof in this report; a Lean formalization is accepted with its single Newton-factorization input visible as a theorem hypothesis. [\[paper proof in this report\]](#) [\[kernel-verified Lean\]](#)

The one-logarithm deficit (Section 22). To make the semigroup determinant nearly force a gap comparable with X , one needs $r \asymp (\log X)^2$ candidate nodes. Conditional on a counterexample, the exact ProveIt classification supplies only $\asymp \log X$ candidates in $[X, 2X]$. Purely archimedean interpolation, at its present strength, therefore misses the conditional population by exactly one power of $\log X$. A successful refinement must either reduce the determinant’s node requirement to $O(\log X)$ or extract additional nonarchimedean size from the same nodes. [\[paper proof in this report\]](#)

Matrix coefficients and sparse interpolation (Section 65). The conjecture is exactly a restricted 2×2 Matrix Coefficient principle. Independently, every nonzero relation supported on m monomials restricts to an exponential polynomial with at most $m - 1$ positive-curve zeros; the matching positive-node interpolant has all $R + 1$ coefficients nonzero. Thus R counterexample points have exact sparse interpolation complexity $R + 1$, so saving even one monomial over ordinary interpolation would rule them out. The matrix equivalence, sparse lower bound, nonzero integral determinant, and matching interpolant are kernel checked; the BK-area calibration and resulting factor-four comparison with the general Matrix Coefficient target remain paper-level. [\[kernel-verified Lean\]](#) [\[paper proof in this report\]](#)

1.4 Why this is not a proof

The conjecture is already a special integral instance of the real four exponentials conjecture. In the natural 2×2 matrix

$$\begin{pmatrix} 1 \cdot \log 2 & 1 \cdot \log 3 \\ x \log 2 & x \log 3 \end{pmatrix},$$

the four exponentials are $2, 3, 2^x, 3^x$. For a nonintegral solution, both the row pair and the column pair are rationally linearly independent, yet all four exponentials are algebraic. The conclusion needed here is therefore exactly what four exponentials predicts and what the known six exponentials theorem does not supply in a 2×2 configuration.

The additional integrality hypotheses are stronger than mere algebraicity, and they do yield the new determinant inequalities of Sections 19 and 20 and the bilinear criterion of Section 12. They do not, by themselves, produce enough nodes, enough lower-bound divisibility, or the required independence of logarithm products to cross the four-exponentials threshold. Conjecture 2.1 therefore remains marked open throughout this report.

2 The problem and its equivalent forms

2.1 Original statement

Conjecture 2.1 (Alaoglu–Erdős, two-base integer form). *For every $x \in \mathbb{R}$,*

$$2^x \in \mathbb{Z} \quad \text{and} \quad 3^x \in \mathbb{Z} \quad \implies \quad x \in \mathbb{Z}.$$

Status: recorded in Lean as the proposition `AlaogluErdosConjecture`, without being asserted.

Since 2^x and 3^x are positive, the hypotheses actually put them in $\mathbb{N}$; and $2^x \in \mathbb{N}$ forces $x \geq 0$, since a negative exponent gives $0 < 2^x < 1$. All solutions therefore live in $[0, \infty)$ (`nonneg_of_two_rpow_integer`), and “integer” may be replaced throughout by “positive natural number.”

Historical placement. Questions of this type go back to the 1944 study of colossally abundant numbers by Alaoglu and Erdős [1], where the arithmetic nature of simultaneous rational powers p^x, q^x of distinct primes controls the quotients of consecutive colossally abundant numbers — a theme reaching back to Ramanujan’s superior highly composite numbers [2]. Alaoglu and Erdős isolated the two-prime assertion explicitly, regarded it as very likely, could not prove it, and reported a communicated three-prime result of Siegel. In modern language Siegel’s case is supplied by the six exponentials theorem, whereas the two-base statement sits at the unresolved 2×2 threshold of the four exponentials conjecture (Schneider’s first problem [9]; see Waldschmidt [10, 13]). The elementary-looking integrality question is thus one of the canonical concrete tests of the gap between six and four exponentials, not an isolated recreational puzzle. As of the dated literature check for this report, the four exponentials conjecture remains open.

2.2 Candidate-base normalization

Put

$$\vartheta := \log_2 3 = \log_2 3 = 1.584962500721156 \dots$$

and define the *candidate set*

$$\mathcal{C} := \{m \in \mathbb{N} : m > 0, m^\vartheta \in \mathbb{N}\}.$$

Lemma 2.2 (Exponent–candidate bijection). *The maps $x \mapsto 2^x$ and $m \mapsto \log_2 m$ are mutually inverse bijections between $\mathcal{S} := \{x \in \mathbb{R} : 2^x, 3^x \in \mathbb{Z}\}$ and $\mathcal{C}$. Under the bijection, $3^x = m^\vartheta$.*

Proof. If $x \in \mathcal{S}$, set $m = 2^x \in \mathbb{N}$. Then $m^\vartheta = \exp(\log m \cdot \log_2 3) = \exp(x \log 3) = 3^x \in \mathbb{N}$. Conversely, if $m \in \mathcal{C}$ and $x = \log_2 m$, then $2^x = m$ and the same identity gives $3^x = m^\vartheta \in \mathbb{N}$. $\square$

Status: [\[kernel-verified Lean\]](#) —

`exists_twoBaseNaturalCandidate_of_two_three_rpow_integer,`
`TwoBaseNaturalCandidate.integer_powers.`

Integral solutions correspond exactly to powers of two, by the identifier

$$\text{twoBaseCandidateExponent_isInteger_iff_isPowerOfTwo},$$

so

$$\boxed{\text{Conjecture 2.1} \iff \mathcal{C} = \{2^n : n \in \mathbb{N}_0\}}. \tag{1}$$

This equivalence is itself formalized, as

$$\text{alaogluErdosConjecture_iff_candidates_are_powers_of_two}.$$

The candidate formulation is well suited to unique factorization, finite certification, finite differences of t^θ , multiplicative-rank arguments, and the determinant constructions of this report, and it is the formulation in which most of the corpus is stated.

Two closure properties of $\mathcal{S}$ are elementary but essential: $\mathcal{S}$ is closed under addition and under multiplication by positive naturals, so it is an additive submonoid of $[0, \infty)$ containing $\mathbb{N}$ (`rpow_integer_add`, `rpow_integer_nat_mul`).

2.3 The logarithmic determinant

Writing

$$M = 2^x, \quad A = 3^x,$$

a solution is equivalently a positive integer pair satisfying

$$\log M \log 3 - \log A \log 2 = 0, \tag{2}$$

and the conjecture says that all positive integer solutions of (2) are

$$(M, A) = (2^n, 3^n), \quad n \in \mathbb{N}_0.$$

Thus the problem is the vanishing of a 2×2 determinant of logarithms of algebraic numbers,

$$\det \begin{pmatrix} \log M & \log 2 \\ \log A & \log 3 \end{pmatrix} = 0,$$

with the additional information that two of the four entries are logarithms of *unknown integers*. This is not merely an analogy with the four exponentials conjecture; it is its exact local geometry.

2.4 One-parameter subgroup form

Let

$$\Gamma = \{(2^t, 3^t) : t \in \mathbb{R}\} \subset \mathbb{R}_{>0}^2.$$

Then Conjecture 2.1 is the statement

$$\Gamma \cap \mathbb{N}^2 = \{(2^n, 3^n) : n \in \mathbb{N}_0\}.$$

The curve Γ is a real one-parameter subgroup of the multiplicative torus $\mathbb{G}_m^2(\mathbb{R})$, and the assertion is that it meets the integer lattice points of the torus only at the obvious places. The difficulty is that its direction vector $(\log 2, \log 3)$ is transcendental rather than algebraic, so standard algebraic-subgroup intersection theorems do not apply directly.

2.5 The four-exponentials reduction

Suppose x is nonintegral and $2^x, 3^x$ are algebraic. The rational-exponent argument (Lemma 2.3) shows that x cannot be rational, so $1, x$ are $\mathbb{Q}$ -linearly independent; and $\log 2, \log 3$ are $\mathbb{Q}$ -linearly independent since $2^a 3^b = 1$ forces $a = b = 0$. The four exponentials built from these two pairs are

$$e^{1 \cdot \log 2} = 2, \quad e^{1 \cdot \log 3} = 3, \quad e^{x \log 2} = 2^x, \quad e^{x \log 3} = 3^x,$$

all four algebraic — contrary to the four exponentials conjecture. This reduction is formalized: `alaogluErdosConjecture_of_realFourExponentialsConjecture`. [\[kernel-verified Lean\]](#)

The reduction is a diagnostic, not just context: any proposed elementary argument must exploit integrality, positivity, or arithmetic structure beyond bare algebraicity of those four values, or else it would prove a substantial case of the four exponentials conjecture by another name.

2.6 Rational and algebraic exponents

Lemma 2.3 (Rational exponent). *If $x \in \mathbb{Q}$ and $2^x \in \mathbb{Z}$, then $x \in \mathbb{Z}$. Consequently every nonintegral solution is irrational.*

Proof. Write $x = a/b$ in lowest terms with $b > 0$. If $2^{a/b} = m \in \mathbb{N}$, then $2^a = m^b$; comparing the exponent of 2 gives $a = b v_2(m)$, so $b \mid a$ and $b = 1$. $\square$

Status: [\[kernel-verified Lean\]](#) —

`IntegerExponent.integer_of_rational_of_two_rpow_integer,`
`irrational_of_not_integer_of_two_rpow_integer.`

Gelfond–Schneider upgrades “irrational” to “transcendental” and removes the algebraic irrational case entirely.

Lemma 2.4 (Transcendence dichotomy). *If $2^x \in \mathbb{Z}$, then either $x \in \mathbb{Z}$ or x is transcendental. In particular every counterexample to Conjecture 2.1 is a transcendental number, and $\vartheta = \log_2 3$ is itself transcendental.*

Status: [\[kernel-verified Lean\]](#) —

`transcendental_of_not_integer_of_two_three_rpow_integer,`
`transcendental_logThreeDivLogTwo;` the supporting Gelfond–Schneider material is discussed in Section 4.

Together with the finite verification these lemmas already remove several whole regimes. In summary form, the pinned corpus proves the following unconditional restrictions on a hypothetical counterexample x .

- x is irrational, and in fact transcendental. [\[kernel-verified Lean\]](#)
- ϑ is transcendental. [\[kernel-verified Lean\]](#)
- $x \geq 12$, by the kernel-checked zero-free region $2^x < 4096$. [\[kernel-verified Lean\]](#)
- $2^x \geq 2^{27}$, hence $x \geq 27$, by the directed-rounding scan; this raises the barrier only under the stated software trust boundary and is not a kernel certificate. [\[software computation\]](#)
- The real four exponentials conjecture implies Conjecture 2.1. [\[kernel-verified Lean\]](#)

These facts eliminate the rational, algebraic-irrational, small, and routine three-exponentials cases. They also explain why an elementary manipulation of logarithms is unlikely to finish the proof: the unknown lies exactly where present transcendence theory changes from theorem to conjecture.

2.7 The three-base theorem and its determinant architecture

By contrast with the two-base problem, three multiplicatively independent bases fall to the six exponentials mechanism, and the special case needed here is fully formalized.

Theorem 2.5 (Three-base theorem). *If $2^x, 3^x, 5^x \in \mathbb{Z}$ for a real x , then $x \in \mathbb{Z}$.*

Status: [\[kernel-verified Lean\]](#) — modules `IntegerExponent`, `IntegerExponent.DeterminantBound`, `IntegerExponent.ExponentialKernel`, `IntegerExponent.Grid`, following Waldschmidt’s square-determinant proof [12].

In outline, for irrational x one forms a large square matrix $\mathcal{A}_{ij} = \exp(a_i b_j)$, where the row arguments a_i encode monomials in $2, 3, 5$ and the column arguments b_j encode affine forms in $1, x$. Irrationality makes the column parameters pairwise distinct and multiplicative independence makes the row parameters pairwise distinct; a total-positivity argument gives $\det \mathcal{A} \neq 0$; integrality of all entries forces $|\det \mathcal{A}| \geq 1$; and an interpolation estimate with explicit size choices gives $|\det \mathcal{A}| < 1$, a contradiction.

The dimension count is exactly where the argument stops short of Conjecture 2.1. With three independent row directions and two column directions the analytic upper bound decays faster than the arithmetic lower bound 1 permits, and the contradiction closes. Shrinking to two bases removes one row direction, and the same balance then lands precisely on the classical boundary where the decay below one fails. Nothing in the proof is wasteful enough to be recovered by sharper constants: a successful two-base determinant of this shape would be a new four-exponentials argument, which is why Section 19 and Section 20 enlarge the *node set* rather than tune the estimate.

A denominator-controlled rational third output. The rational-output extension now formalized is deliberately conditional. Let $a > 1$ and assume the monomial map

$$(i, j, k) \mapsto 2^i 3^j a^k$$

is injective. If $2^x = m_2$ and $3^x = m_3$ are natural numbers, $a^x = q \in \mathbb{Q}$, and for some $u, v \in \mathbb{N}_0$ the reduced denominator satisfies

$$\text{den}(q) \mid m_2^u m_3^v,$$

then $x \in \mathbb{Q}$, and in fact $x \in \mathbb{Z}$ because 2^x is integral. A concrete corollary obtains the monomial injectivity when a has a prime divisor different from 2 and 3.

On the determinant side, a nonzero $m \times m$ real determinant whose entries become integers after multiplication by a common positive denominator Q satisfies $|\det A| \geq Q^{-m}$. The explicit analytic estimate also retains a denominator margin: for $m = 2r$, $r > 2$, $192C \leq r$, and $bs + 2r < r^2$, the verified upper-bound expression remains strictly below 1 even after multiplication by b^s .

Status: [\[kernel-verified Lean\]](#) — `one_div_pow_le_abs_det_of_common_denominator`, `exponential_det_numeric_bound_with_denominator`, `rational_of_two_three_a_rpow_rational_of_den_dvd_outputs`, and `integer_of_two_three_a_rpow_rational_of_prime_factor_of_den_dvd_outputs` in `RationalThirdBase`. These are denominator-divisibility and explicit-margin results; the unrestricted determinant grid for an arbitrary rational third-base output has *not* been completed.

3 Repository timeline and verification status

The formal development advanced in three waves of Lean work, interleaved with four documentation milestones. All commits are in the repository [16]; the mathematical role of each is summarized below.

Commit	Role
f492c18f1	Three-base theorem via interpolation determinants (<code>IntegerExponent</code> and sub-modules).
5ed4e6185	Merge of the three-base branch.
a37d97f9d	Basic two-base module: the conjecture as a proposition, closure under natural multiples, irrationality of nonintegral solutions, endpoint gaps, numerical bounds for ϑ , the $2^x < 10$ exclusion, second-difference obstruction, aligned-block 2/3 counting.
4a9b03d12	Major corpus: finite checks through $2^x < 256$; arithmetic descent; multiplicative rank; candidate structure; zero density; shrinking targets; the 3-smooth theorem; the four-exponentials reduction; transcendence consequences; and a complete Gelfond–Schneider port.
18a85e6b8	Merge of the two-base branch.
423ac24f0	First survey document (<code>two-base-partial-results.tex</code>) with rendered PDF.
98646b455	Three further modules, kernel-verified: <code>FractionalPartDensity</code> (density dichotomy and gap equivalences), <code>OddCore</code> (common odd core, improved counting, dyadic blocks), <code>FiniteCheck4096</code> (zero-free region to $2^x < 4096$).
716617d72	Survey finalized after kernel acceptance of the 4096 region.
11ac3093f	First independent research report (audit at 423ac24f0, exact conditional structure, interval scan to 2^{20}); one of the sources merged into the present document.
83ef82ab1	Merge of the two earlier documents into a single expanded report.
46f5b67e3	Localization, sharp arithmetic progressions and an exact generator: <code>Localization</code> , <code>SharpProgression</code> , and the first form of <code>PrimitiveGenerator</code> .
5d3982ceb	Merge of the parallel conditional-structure and asymptotic formalizations: canonical primitive generator and exact solution monoid, least solution, exact dyadic counting, the quadratic cumulative asymptotic, block Weyl sums and block equidistribution, simultaneous output normalization and power-of-three denominator normalization, the rational-power index, the radical modules (radical reformulation, primitive radical, exact radical degree), the higher-difference obstruction, and denominator-controlled rational third-base outputs.
8c9c59224	Denominator-growth constraint $q_{n+1} < 2m^{q_n}$ on the continued fraction of a hypothetical solution (<code>DenominatorGrowth</code>), together with the report analysis naming the obstruction as a vanishing integer-coefficient linear form in products of two prime logarithms.

Commit	Role
2ece309cc	Second research report: arbitrary-node Vandermonde repulsion, the semigroup generalized-Vandermonde hierarchy and its dyadic $O((\log X)^2)$ bound, the one-logarithm deficit analysis, and the compact-orbit and block-growth reformulations; the other source merged into the present document.

The audit performed in 11ac3093f treated results of the then-uncommitted modules as article-level claims; with 98646b455 in the tree and kernel-verified, those results are now cited as committed Lean. The interval scan to 2^{20} remains at the [\[software computation\]](#) level by design. The feature modules merged in 5d3982ceb prove the cumulative quadratic asymptotic and the finite-trigonometric-polynomial block averages; they do *not* yet formalize the stronger continuous-test-function or interval-counting equidistribution theorem, and that distinction is maintained throughout this report.

3.1 Evidence convention

The evidence convention of this report is source-tree plus build-level, not rebuild-on-demand. Every claim marked [\[kernel-verified Lean\]](#) was checked against the cited Lean source and theorem statement at the repository state above, and the corresponding module is on the aggregate import surface `ExponentialIdentities.lean`, where it is recorded as accepted by the Lean kernel. The label does not assert that a separate fresh build was run while preparing this document; it asserts that the theorem exists, in the stated form, on an import surface the kernel has accepted.

Three weaker labels are used deliberately. [\[software computation\]](#) marks a finite interval computation with an explicit software trust boundary and no kernel certificate; [\[paper proof in this report\]](#) marks a conventional proof given in this report, written with formalization in mind but not yet translated; and [\[Lean draft, not yet accepted\]](#) marks a Lean module that exists in the tree as a draft but has *not* been accepted by the kernel, and therefore carries no more weight than [\[paper proof in this report\]](#). The complete label-by-label inventory is in Appendix A, and the per-module inventory in Appendix B.

4 The kernel-verified corpus

Throughout this section x denotes a hypothetical nonintegral solution, $m = 2^x \in \mathbb{N}$, $n = 3^x \in \mathbb{N}$. Every statement here is [\[kernel-verified Lean\]](#) unless labeled otherwise.

4.1 Fractional-part gaps and shrinking targets

An integral power caught strictly between consecutive integral powers of its base must clear both by at least one full unit; dividing through normalizes this to the fractional part.

Theorem 4.1 (Fractional-part gaps). *Let $x \in (N, N + 1)$ with $N \in \mathbb{N}$, $t = x - N$, and suppose $2^x, 3^x \in \mathbb{Z}$. Then*

$$1 + 2^{-N} \leq 2^t \leq 2 - 2^{-N}, \quad 1 + 3^{-N} \leq 3^t \leq 3 - 3^{-N}.$$

Status: [\[kernel-verified Lean\]](#) — `two_three_fractional_part_gaps`,
`two_three_rpow_gaps_near_nat`, from
`rpow_integer_gap_between_consecutive_powers`.

Near the endpoints the excluded zones have width $\asymp 2^{-N}$: they *shrink* as the integer part grows. Applied to all multiples kx (which are again solutions), this yields explicit Diophantine lower bounds.

Theorem 4.2 (Effective irrationality bound). *Let $x \notin \mathbb{Z}$ satisfy $2^x \in \mathbb{Z}$ and write $m = 2^x$. Then for all integers p and all $k \geq 1$,*

$$\left| x - \frac{p}{k} \right| \geq \frac{1}{2 k m^k}.$$

Exact (slightly stronger) radii, and two-base versions taking the better of the base-2 and base-3 radii, are formalized as well.

Status: [\[kernel-verified Lean\]](#) — `rational_approximation_exponential_lower_bound`,
`two_three_multiple_distance_to_integer_lower_bound`,
`exists_natural_base_rational_approximation_lower_bound` in `ShrinkingTarget`.

In shorthand, $\|kx\|_{\mathbb{Z}} \gtrsim 2^{-kx}$. This bound is compatible with ordinary Diophantine approximation: Dirichlet guarantees infinitely many k with $\|kx\|_{\mathbb{Z}} < 1/k$, and the shrinking radius 2^{-kx} is incomparably smaller. The interplay between this exponential gap and equidistribution is discussed in Section 26.

The bound dualizes into a growth constraint on the continued fraction of a solution: if $|x - p/q| < 1/(qQ)$ with $q, Q \geq 1$, then $Q < 2m^q$. Applied to consecutive convergent denominators, which satisfy $|x - p_n/q_n| < 1/(q_n q_{n+1})$, this gives

$$q_{n+1} < 2m^{q_n} :$$

the partial quotients of a hypothetical nonintegral solution are capped by a single exponential in the preceding denominator, with base its own output $m = 2^x$.

Status: [\[kernel-verified Lean\]](#) —
`approximation_quality_lt_of_noninteger_solution` in `DenominatorGrowth`.

4.2 Verified zero-free regions and certificate architecture

Theorem 4.3 (Zero-free region). *If $2^x, 3^x \in \mathbb{Z}$ and $2^x < 4096$, then $x \in \mathbb{Z}$. Equivalently, every nonintegral solution satisfies $x \geq 12$, and every candidate that is not a power of two is at least 4096.*

Status: [\[kernel-verified Lean\]](#) —
`integer_of_two_three_rpow_integer_of_two_rpow_lt_4096`,
`twelve_le_of_not_integer_of_two_three_rpow_integer`,
`le_of_twoBaseNonintegerCandidate_4096` in `FiniteCheck4096`, extending `FiniteCheck`
 $(2^x < 100)$ and `FiniteCheck256` ($2^x < 256$, hence $x \geq 8$).

The certificate architecture is worth recording, both because it is the repository’s model of a trustworthy large computation and because scaling it is a named research direction. For each non-power-of-two m the certificate is a pair of exact integer power comparisons against rational enclosures $L_p/L_q < \vartheta < U_p/U_q$ drawn from continued-fraction convergents of ϑ : from

$$a^{L_q} < m^{L_p} \quad \text{and} \quad m^{U_p} < (a+1)^{U_q}$$

one traps $a < m^\vartheta < a+1$, so $m^\vartheta \notin \mathbb{Z}$. The $2^x < 256$ stage verifies per-value certificates by `norm_num` with three enclosure tiers (exponent sizes ~ 500 , ~ 1500 , ~ 25000). The extension to 4096 verifies 3836 certificates through a single kernel-replayed `decide` over a table generated by PARI/GP and revalidated independently with exact big-integer arithmetic in Python. Five tiers suffice, with tier populations 78, 488, 3241, 28, 1:

tier	lower convergent L_p/L_q	upper convergent U_p/U_q
0	569/359	485/306
1	1054/665	1539/971
2	25781/16266	24727/15601
3	50508/31867	125743/79335
4	176251/111202	301994/190537

Only the value $m = 2762$, whose ϑ -power is within $3 \cdot 10^{-9}$ of an integer in logarithmic scale, needs tier 4; its certificate compares integers of roughly 90,000 digits, which the kernel’s exact arithmetic replays without difficulty. No axioms beyond the Lean kernel are involved; in particular no `native_decide`.

4.3 Arithmetic descent

Theorem 4.4 (Affine descent). *If the two integral outputs factor simultaneously as $2^x = 2^r u^k$ and $3^x = 3^r v^k$ with $r \geq 0$, $k \geq 1$, then $(x - r)/k$ is again a solution. Combined with the $x \geq 8$ region this gives the unconditional constraint*

$$r + 8k \leq x$$

for nonintegral x ; in particular common perfect-power shapes of the two outputs are strongly restricted.

Status: [\[kernel-verified Lean\]](#) — `affine_descent_integral_powers`, `base_depth_add_eight_mul_degree_le`, `min_output_base_factorization_add_eight_le` in `ArithmeticRigidity`. With the 4096 region the constant 8 improves to 12 by bookkeeping; with the interval scan of Section 25 it would improve to 20 at the [\[software computation\]](#) level.

For the *least* counterexample the same descent later becomes an exact primitivity statement in Section 5: no nontrivial matched decomposition exists at all.

4.4 Local progression obstructions and zero density

The function $t \mapsto t^\vartheta$ is strictly convex with second differences at moderate steps lying strictly between 0 and 1 — hence never integers.

Theorem 4.5 (Committed AP obstruction). *Let $n, h \geq 1$ with $h^5 \leq n$. Then $n, n + h, n + 2h$ are not all candidates. In particular no three consecutive naturals are candidates, and among the positive integers up to $3N$ at most $2N$ are candidates.*

Status: [\[kernel-verified Lean\]](#) —

`not_three_arithmetic_progression_twoBaseNaturalCandidates`,
`not_three_consecutive_logThreeDivLogTwo_rpow_integers`,
`twoBaseNaturalCandidate_Icc_card_le`. The sharper criterion $\vartheta(\vartheta - 1)h^2 < n^{2-\vartheta}$ is
kernel-verified in `SharperProgressions`, and the analogous four-term criterion is
kernel-verified in `ThirdDifference`. The arbitrary-order statement for every $r \geq 2$ is
`not_arithmetic_progression_twoBaseNaturalCandidates_of_curvature` in
`HigherDifference`, which formalizes a generic iterated forward-difference mean-value
identity and applies uniformly in the order r .

For finite verification the same curvature mechanism is also available in a purely rational effective form: if $h^{400} \leq n^{83}$ — step exponent $83/400 = 0.2075$, against the $1/5$ of Theorem 4.5 — then $n, n + h, n + 2h$ are not all candidates.

Status: [\[kernel-verified Lean\]](#) —

`not_three_arithmetic_progression_twoBaseNaturalCandidates_sharp` in
`SharpProgression`, with the kernel-replayed enclosure $\vartheta < 317/200$ obtained from the exact
integer comparison $3^{200} < 2^{317}$.

The hypothesis bounding h is essential: for steps comparable to n the second difference is large and the obstruction vanishes. This locality is a genuine barrier — a *global* progression obstruction would, for example, forbid candidates with odd part $2^s \pm 1$ via progressions such as $(2^i, 2^{i+t-1}, 2^i(2^t - 1))$, but no such global statement is available; under a counterexample the exact monoid of Section 5 shows sparse candidates whose additive relations the curvature argument cannot see.

Feeding the local obstruction into the theory of Roth numbers gives quantitative sparsity.

Theorem 4.6 (Zero density). *For every $H \geq 1$ the number of candidates below N is at most $H^5 + \lceil (N - H^5)/H \rceil r_3(H)$, where $r_3(H)$ is the maximal size of a 3-AP-free subset of $\{0, \dots, H - 1\}$. Consequently the candidate counting function is $o(N)$: the candidates, and a fortiori the values 2^x of nonintegral solutions, have asymptotic density zero.*

Status: [\[kernel-verified Lean\]](#) — `twoBaseNaturalCandidatesBelow_card_le`,
`twoBaseNaturalCandidatesBelow_isLittleO_id`,
`tendsto_twoBaseNonintegerCandidatesBelow_density_zero` in `ZeroDensity` and
`NonintegerDensity`.

4.5 Multiplicative rank two and the common odd core

Theorem 4.7 (Rank two). *The prime-exponent vectors $\mathbf{v}(a) = (v_p(a))_p$ of all candidates span a rational space of dimension at most two. Consequently at most $(\log_2 N)^2$ candidates lie below N .*

Status: [\[kernel-verified Lean\]](#) — `finrank_span_primeExponentVectors_le_two`, `twoBaseNaturalCandidatesBelow_card_le_clog_sq` in `MultiplicativeRank`. The mechanism abstracts the determinant method: three candidates with independent exponent vectors would give three multiplicatively independent bases, contradicting the three-base theorem of Section 2 through the irrationality of ϑ (`irrational_logThreeDivLogTwo`).

Since 2 is itself a candidate, all odd parts lie on at most one further rational direction. Pairwise this yields the cross-valuation identity, and globally a common odd core.

Lemma 4.8 (Cross-valuation identity). *Let $a, b \in \mathcal{C}$ and let p be an odd prime with $v_p(a) > 0$. Then for every odd prime q ,*

$$v_p(a) v_q(b) = v_p(b) v_q(a).$$

Status: [\[kernel-verified Lean\]](#) — `odd_factorization_cross_mul` in `CandidateStructure`; the fixed-power relations `fixed_power_relation_of_odd_prime_factor` follow.

Theorem 4.9 (Common odd core). *There is a single odd $w > 1$ such that every candidate equals $2^i w^j$ for some $i, j \in \mathbb{N}_0$; candidates that are not powers of two have $j \geq 1$. In exponent coordinates: every solution has the form $x = i + j \log_2 w$ with $i, j \in \mathbb{N}_0$.*

Status: [\[kernel-verified Lean\]](#) — `exists_common_odd_core`, `exists_common_odd_core_split`, `exists_common_odd_core_exponent` in `OddCore`. The canonical (gcd-normalized) version with uniqueness of w and of each coordinate pair is `existsUnique_primitive_common_odd_core` in `PrimitiveOddCore`.

Corollary 4.10 (Improved counting). *The number of candidates below N is at most $\log_2 N \cdot \log_3 N$, and each dyadic block $[2^T, 2^{T+1})$ contains at most $\log_3(2^{T+1}) \approx 0.631(T+1)$ candidates. In exponent coordinates, the number of solutions with $x \in [T, T+1)$ grows at most linearly in T , against the quadratic global bound.*

Status: [\[kernel-verified Lean\]](#) — `twoBaseNaturalCandidatesBelow_card_le_clog_mul_clog`, `twoBaseNonintegerCandidatesBelow_card_le_clog_mul_clog`, `twoBaseNaturalCandidatesInDyadicBlock_card_le` in `OddCore`.

The logarithmic-square order resists improvement by counting alone: if a candidate with an odd prime factor exists, multiplicative closure generates injective boxes $2^i a^j$, so the count below $2^k a^l$ is at least kl (`twoBaseNaturalCandidatesBelow_card_ge_box`, `twoBaseNonintegerCandidatesBelow_card_ge_box`) — a conditional lower bound of the same quadratic-in-logarithm shape, made exact in Section 5.

4.6 Transcendence

The corpus contains a complete Lean proof of the Gelfond–Schneider theorem (Hilbert’s seventh problem), following Hua’s exposition [4], ported from Michail Karatarakis’s Mathlib formalization (Mathlib PR #42911, Apache 2.0).

Theorem 4.11 (Gelfond–Schneider). *If α, β are algebraic, $\alpha \neq 0, 1$, and β is irrational, then α^β is transcendental.*

Status: [\[kernel-verified Lean\]](#) —

GelfondSchneider.transcendental_cpow_of_isAlgebraic_of_irrational.

Corollary 4.12 (Algebraic–integral dichotomy). *If $2^x \in \mathbb{Z}$ then x is either an integer or transcendental; algebraicity of such x is equivalent to integrality. Every nonintegral two-base solution is transcendental, and $\vartheta = \log_2 3$ is transcendental.*

Status: [\[kernel-verified Lean\]](#) —

transcendental_of_not_integer_of_two_rpow_integer,
isAlgebraic_iff_integer_of_two_rpow_integer,
transcendental_of_not_integer_of_two_three_rpow_integer,
transcendental_logThreeDivLogTwo.

The dichotomy does not resolve the conjecture — the unknown exponent is allowed to be transcendental — but it constrains every construction below: the primitive generator β of Section 5 is transcendental, as is $\log_2 w$ for the odd core.

4.7 Classification of auxiliary bases: the 3-smooth theorem

Theorem 4.13 (3-smooth classification). *Suppose $x \notin \mathbb{Z}$ and $2^x, 3^x \in \mathbb{Z}$. Then for a positive natural a , the power a^x is an integer if and only if $a = 2^u 3^v$ is 3-smooth. In particular a third base with any prime factor other than 2 and 3 forces $x \in \mathbb{Z}$.*

Status: [\[kernel-verified Lean\]](#) —

rpow_integer_iff_eq_two_pow_mul_three_pow_of_not_integer,
integer_of_two_three_a_rpow_integer_of_prime_factor in ThreeSmooth.

This is an exact classification of bases for a *fixed* hypothetical counterexample. It should not be confused with the classification of candidate values 2^x as x ranges over all solutions, which is the subject of the next section. Note also the scope limitation explained in Section 26: the theorem controls the quantity b^x , not divisibility of the integer 2^x by b .

4.8 Unconditional gap equivalences

Because natural multiples of a solution are solutions, one nonintegral solution generates infinitely many, with irrational generator; Dirichlet approximation and a rotation-stepping argument spread their fractional parts densely.

Theorem 4.14 (Density dichotomy). *If some $x \notin \mathbb{Z}$ satisfies $2^x, 3^x \in \mathbb{Z}$, then for every $t \in [0, 1]$ and $\varepsilon > 0$ there is a nonintegral solution y with $|\{y\} - t| < \varepsilon$; and the set of nonintegral solutions is infinite.*

Theorem 4.15 (Gap equivalences). *Each of the following is equivalent to Conjecture 2.1:*

- (i) *some nonempty open interval $(u, v) \subseteq (0, 1)$ contains no fractional part of a nonintegral solution;*
- (ii) *the fractional parts of nonintegral solutions are bounded away from 1;*
- (iii) *the fractional parts of nonintegral solutions are bounded away from 0.*

Status: [\[kernel-verified Lean\]](#) — `dense_fract_of_noninteger_solution`,
`infinite_noninteger_solutions_of_exists`,
`alaogluErdosConjecture_iff_fract_avoids_interval`,
`alaogluErdosConjecture_iff_fract_bounded_away_one`,
`alaogluErdosConjecture_iff_fract_bounded_away_zero` in `FractionalPartDensity`,
built on the reusable `exists_pos_nat_fract_close_of_irrational`.

Remark 4.16 (Consistency pincer). Theorem 4.1 bounds $\{x\}$ away from 0 and 1 by margins $\asymp 2^{-\lfloor x \rfloor}$ that decay with the height of the solution, while Theorem 4.14 produces fractional parts arbitrarily close to 0 and 1 along multiples of unbounded height. The two results meet edge-to-edge without contradiction, and Theorem 4.15 shows the decay is essential: *any* height-independent gap, however small and wherever located, is exactly as strong as the full conjecture. The elementary gap bounds are therefore of optimal shape, and the fractional-part axis measures the entire remaining distance to Conjecture 2.1.

These equivalences are *unconditional* biconditionals. Conditional on failure, the exact structure of the next section gives kernel-verified convergence for every finite trigonometric polynomial and the paper-level strengthening to full blockwise Weyl equidistribution (Section 5).

5 Exact conditional structure of all solutions

The most consequential fact in the current corpus is not another exclusion range. It is the complete description of *all* solutions under the hypothesis that even one nonintegral solution exists. Conditional on failure of Conjecture 2.1, the solution set is not merely sparse or contained in some rank-two set: it is a free rank-two additive monoid with an explicit canonical generator, exactly countable in every dyadic block, and equidistributed blockwise on the circle.

Throughout this section we write

$$\mathcal{S} := \{x \in \mathbb{R} : 2^x \in \mathbb{N} \text{ and } 3^x \in \mathbb{N}\}, \quad \mathcal{C} := \{2^x : x \in \mathcal{S}\}, \quad \mathcal{P} := \{(2^x, 3^x) : x \in \mathcal{S}\},$$

and we say a counterexample exists when $\mathcal{S} \not\subseteq \mathbb{N}_0$. The canonical primitive odd core, the rational-power index, the power-of-three denominator normalization, the unique least solution, the simultaneous primitive output normalization, the exact free-monoid classification, the exact unit-block and dyadic counts, the cumulative quadratic asymptotic, and the finite-Fourier block averages are all [\[kernel-verified Lean\]](#). Only the extension of the block averages from finite trigonometric polynomials to arbitrary continuous tests or interval indicators (Theorem 5.10) remains [\[paper proof in this report\]](#).

5.1 The primitive common odd core

Theorem 5.1 (Primitive common odd core). *Assume $\mathcal{C}$ contains a non-power of two. Then there is a unique odd integer $w > 1$ that is power-free in the sense*

$$\gcd\{v_p(w) : p \mid w\} = 1,$$

and every candidate $b \in \mathcal{C}$ has a unique representation

$$b = 2^i w^j, \quad i, j \in \mathbb{N}_0.$$

Proof. Choose a non-power-of-two candidate a . Its odd valuation vector $A = (A_q)_{q \neq 2}$, $A_q = v_q(a)$, is nonzero and finitely supported. Let

$$g = \gcd\{A_q : A_q > 0\}, \quad c_q = A_q/g, \quad w = \prod_{q \neq 2} q^{c_q}.$$

Then w is odd, greater than one, and the positive c_q have gcd one.

Let $b \in \mathcal{C}$ and set $B_q = v_q(b)$. Fix an odd prime p with $A_p > 0$. By the cross-valuation identity of Section 4, $A_p B_q = B_p A_q$ for all odd q . If $A_q = 0$, this forces $B_q = 0$. On the support of A it says

$$B_q = \frac{B_p}{A_p} A_q = \left(\frac{g B_p}{A_p} \right) c_q.$$

Put $t = g B_p / A_p \in \mathbb{Q}_{\geq 0}$. Every $t c_q$ is an integer. Since the c_q have gcd one, there are integers r_q , finitely many nonzero, with $\sum r_q c_q = 1$; therefore $t = \sum r_q (t c_q) \in \mathbb{Z}$. Writing $j = t \in \mathbb{N}_0$, the odd part of b is exactly w^j , so $b = 2^i w^j$ with $i = v_2(b)$.

Uniqueness of (i, j) for fixed w follows from the odd valuations and the 2-adic valuation. If w' is another primitive common core, the odd valuation vectors of w and w' generate the same rational ray and are both primitive integer vectors, hence equal. $\square$

Status: [\[kernel-verified Lean\]](#)

`existsUnique_primitive_common_odd_core` in `PrimitiveOddCore`. The module defines `oddPrimeValuationGCD`, proves its equivalence with power-theoretic primitivity for $w > 1$, and kernel-checks both uniqueness of the normalized core and uniqueness of every pair (i, j) . The earlier non-normalized existence statement `exists_common_odd_core` recorded in Section 4 is subsumed by it.

The gcd normalization is exactly what makes w canonical: it prevents a hidden perfect power from being absorbed into the exponent j , so that both the core and the coordinates are determined by $\mathcal{C}$ alone.

Set

$$\alpha := \log_2 w, \quad E := 3^\alpha = w^\vartheta.$$

Every candidate exponent then has the form $x = i + j\alpha$. The number α is irrational (otherwise the rational-exponent lemma applied to $2^\alpha = w$ would make it integral, forcing the odd $w > 1$ to be a power of two), and in fact transcendental by the algebraic–integral dichotomy of Section 4. The candidate condition becomes

$$3^{i+j\alpha} = 3^i E^j \in \mathbb{N},$$

and at this stage the admissible pairs (i, j) still form an unknown subset of $\mathbb{N}_0^2$. The next subsection shows that this subset is completely rigid.

5.2 The rational-power index and the denominator normalization

Because a non-power-of-two candidate exists, some $j > 0$ has $E^j \in \mathbb{Q}$. Let

$$d := \min\{j \in \mathbb{N}_{>0} : E^j \in \mathbb{Q}\}.$$

Lemma 5.2 (All rational powers are multiples of the least one). *For every $j \in \mathbb{Z}$: $E^j \in \mathbb{Q} \iff d \mid j$.*

Proof. The set $J := \{j \in \mathbb{Z} : E^j \in \mathbb{Q}\}$ is an additive subgroup of $\mathbb{Z}$ (it contains 0, products add exponents, and reciprocals negate them), and every nonzero subgroup of $\mathbb{Z}$ equals $d\mathbb{Z}$ for its least positive element d . $\square$

Status: [\[kernel-verified Lean\]](#)

`rationalPowerIndices`, `AddSubgroup.Int.exists_least_positive_generator`, and `exists_rationalPowerIndex` in `RationalPowerIndex`.

Write the positive rational E^d in lowest terms. Its denominator has no prime other than 3: taking a candidate with odd-core exponent $j = dk > 0$ and 2-exponent i , integrality of $3^i(E^d)^k$ shows that any other prime in the reduced denominator would survive in the k -th power. Hence there are unique integers $a \geq 0$ and $A \geq 1$ with

$$E^d = \frac{A}{3^a}, \quad \text{and} \quad 3 \nmid A \text{ when } a > 0, \quad (3)$$

so that $a = 0$ or $3 \nmid A$ in every case. Define

$$\beta := a + d\alpha = a + d \log_2 w, \quad M := 2^\beta = 2^a w^d, \quad A = 3^\beta = 3^a E^d. \quad (4)$$

Thus β is itself a solution, and it is irrational because α is. (Where the denominator normalization is written with a named integer numerator c , the identity $E^d = c/3^a$ with $a = 0$ or $3 \nmid c$ is exactly (3) with $c = A = 3^\beta$.)

Status: [\[kernel-verified Lean\]](#)

The denominator-support step is `exists_three_pow_denominator_of_rational_power` in `ThreeDenominatorNormalization`; the normalized integrality criterion is `three_pow_mul_normalized_pow_integer_iff` in `RationalPowerIndex`.

Equation (3) also has a field-theoretic reading: E is a root of the binomial $X^d - A/3^a$, which is the minimal polynomial of E over $\mathbb{Q}$, so $[\mathbb{Q}(E) : \mathbb{Q}] = d$. That computation, together with the equivalent one-number and localized-radical formulations of the conjecture that follow from it, is developed in Section 6; we do not repeat it here.

5.3 The exact primitive-generator theorem

Theorem 5.3 (Exact primitive generator and solution monoid). *Assume Conjecture 2.1 is false. With w, d, a, β, M, A as above:*

(i) every solution has a unique representation $x = n + k\beta$ with $n, k \in \mathbb{N}_0$, that is,

$$\mathcal{S} = \{n + k\beta : n, k \in \mathbb{N}_0\};$$

(ii) every candidate has a unique representation $m = 2^n M^k$ with $n, k \in \mathbb{N}_0$, that is,

$$\mathcal{C} = \{2^n M^k : n, k \in \mathbb{N}_0\};$$

(iii) every output pair has a unique representation $(2^x, 3^x) = (2^n M^k, 3^n A^k)$;

(iv) β is the unique least nonintegral solution.

Proof. Let $x \in \mathcal{S}$. By Theorem 5.1, $2^x = 2^i w^j$ for unique $i, j \in \mathbb{N}_0$, so $x = i + j\alpha$. From $3^x = 3^i E^j \in \mathbb{N}$ we get $E^j \in \mathbb{Q}$, so Lemma 5.2 gives $j = dk$ for some $k \in \mathbb{N}_0$. Using (3),

$$3^x = 3^i (E^d)^k = 3^{i-ak} A^k.$$

Because $A/3^a$ is reduced, this is an integer exactly when $i \geq ak$. Put $n = i - ak \in \mathbb{N}_0$; then

$$x = i + dk\alpha = n + k(a + d\alpha) = n + k\beta.$$

Conversely, for any $n, k \in \mathbb{N}_0$,

$$2^{n+k\beta} = 2^n M^k \in \mathbb{N}, \quad 3^{n+k\beta} = 3^n A^k \in \mathbb{N},$$

so $n + k\beta \in \mathcal{S}$. If $n + k\beta = n' + k'\beta$ with $k \neq k'$, then $\beta = (n' - n)/(k - k') \in \mathbb{Q}$, a contradiction; hence the representation is unique. This proves (i), and (ii)–(iii) follow by exponentiation. A nonintegral solution has $k \geq 1$; the least such value is attained at $(n, k) = (0, 1)$ and equals β , uniquely so by uniqueness of the representation. $\square$

Status: [\[kernel-verified Lean\]](#)

`exists_canonical_primitiveGenerator_of_not_alaogluErdosConjecture` in `PrimitiveGenerator`, built from

`exists_exact_solution_monoid_of_fixed_common_oddCore` in `RationalPowerIndex`.

The Lean statement supplies the canonical odd core w , the rational-power index d , the depth a , the numerator A , and the unique-coordinate clauses in one package. Conditional on failure, $\mathcal{S}$ is not merely contained in a rank-two additive set: every solution has a unique pair of coordinates in the free additive monoid on 1 and the irrational least generator β . The set of occurring odd-core exponents is exactly $d\mathbb{N}_0$, with the power of two governed by one linear threshold.

Packaged algebraically, Theorem 5.3 says that

$$(\mathbb{N}_0^2, +) \xrightarrow{\sim} (\mathcal{S}, +), \quad (n, k) \mapsto n + k\beta, \quad (\mathbb{N}_0^2, +) \xrightarrow{\sim} (\mathcal{C}, \cdot), \quad (n, k) \mapsto 2^n M^k,$$

are monoid isomorphisms, and that the output-pair monoid $\mathcal{P}$ is freely generated by $(2, 3)$ and (M, A) . Every multiplicative relation among candidates is the obvious coordinate relation; there is no room for an unexpected coincidence. This exact conditional model is the benchmark against which any future density, progression, or descent argument must be tested.

The phrase “least nonintegral” is materially stronger than choosing an arbitrary counterexample. It forces affine primitivity: no nontrivial common root can be extracted from the two output integers while preserving the same base-adic offset. It is the arithmetic analogue of choosing a primitive lattice direction, and the next subsection turns it into exact exclusions.

5.4 Primitivity of the least output pair

Corollary 5.4 (No matched base depth). *For the primitive pair $(M, A) = (2^\beta, 3^\beta)$ we have $\min\{v_2(M), v_3(A)\} = 0$; that is, M is odd or $3 \nmid A$ (possibly both).*

Proof. If both valuations were positive, then $2^{\beta-1} = M/2$ and $3^{\beta-1} = A/3$ would be integers, so $\beta - 1$ would be a smaller nonintegral solution. $\square$

Corollary 5.5 (No common perfect-power degree). *There is no $k \geq 2$ for which both M and A are perfect k -th powers.*

Proof. If $M = U^k$ and $A = V^k$, then $2^{\beta/k} = U$ and $3^{\beta/k} = V$, so β/k is a smaller nonintegral solution. $\square$

Corollary 5.6 (Exact affine primitivity). *If $M = 2^r U^k$ and $A = 3^r V^k$ with $r \in \mathbb{N}_0$, $k \in \mathbb{N}_{>0}$, and $U, V \in \mathbb{N}$, then $r = 0$ and $k = 1$.*

Proof. The affine descent theorem of Section 4 produces the solution $(\beta - r)/k$, which is nonintegral since β is irrational; if $r > 0$ or $k > 1$ it is strictly smaller than β , contradicting leastness. $\square$

Status: [\[kernel-verified Lean\]](#)

```
exists_least_twoBaseNonintegerSolution,
IsLeastTwoBaseNonintegerSolution.unique,
IsLeastTwoBaseNonintegerSolution.not_both_base_dvd_outputs,
IsLeastTwoBaseNonintegerSolution.no_common_perfect_power, and
IsLeastTwoBaseNonintegerSolution.affine_primitive in LeastSolution.
```

The odd-core normalization of M has a symmetric counterpart on the other output, and the two can be imposed simultaneously.

Theorem 5.7 (Simultaneous primitive output normalization). *If the conjecture fails, its least nonintegral solution β admits simultaneous factorizations*

$$M = 2^\beta = 2^a w^d, \quad A = 3^\beta = 3^b v^e,$$

where $w > 1$ is odd and power-primitive, $v > 1$ is power-primitive with $3 \nmid v$, and $d, e > 0$. Moreover $a = 0$ or $b = 0$; the exponents a and b are the exact base-adic depths; d and e are the gcds of the prime valuations of the corresponding base-free parts; and

$$\gcd(\gcd(d, e), |b - a|) = 1.$$

Since every counterexample satisfies $\beta \geq 12$, these normalizations also obey the size dichotomy

$$w^d \geq 4096 \quad \text{or} \quad v^e \geq 3^{12}.$$

Status: [\[kernel-verified Lean\]](#)

`exists_simultaneous_primitive_output_normalization`, `IsLeastTwoBaseNonintegerSolution.simultaneousCoordinateGCD_eq_one`, and `exists_large_primitiveCore_of_not_alaogluErdosConjecture` in `OutputNormalization`. The `gcd-one` statement simultaneously controls the two outer power degrees and the difference of the two base-adic depths; it is strictly stronger than merely excluding a common perfect-power degree (Corollary 5.5).

These restrictions are strong, but they still admit infinitely many formal integer parameter patterns. The unresolved analytic content is untouched: the pair (M, A) must satisfy

$$\frac{\log M}{\log 2} = \frac{\log A}{\log 3},$$

and no known local valuation argument turns the depth dichotomy $a = 0$ or $b = 0$, or the `gcd-one` condition, into a contradiction. Every constraint proved so far is compatible with the existence of such a pair; what fails to exist must fail for a transcendence reason, not a valuation-theoretic one.

5.5 Exact counting

The free-monoid description makes the population of a hypothetical counterexample completely explicit, block by block.

Theorem 5.8 (Exact unit-block and dyadic counts). *Assume a counterexample exists and let β be the primitive generator. For every $N \in \mathbb{N}_0$, the nonintegral solutions in $[N, N + 1)$ are precisely*

$$x_{N,k} = \lceil N - k\beta \rceil + k\beta, \quad 1 \leq k \leq \left\lfloor \frac{N+1}{\beta} \right\rfloor,$$

so their number is exactly $\lfloor (N+1)/\beta \rfloor$. Consequently

$$\#(\mathcal{S} \cap [N, N+1)) = \#(\mathcal{C} \cap [2^N, 2^{N+1})) = 1 + \left\lfloor \frac{N+1}{\beta} \right\rfloor.$$

Proof. For fixed $k \geq 1$ there is at most one integer n with $N \leq n + k\beta < N + 1$; such an $n \geq 0$ exists if and only if $k\beta < N + 1$, and then $n = \lceil N - k\beta \rceil$. Irrationality of β makes $(N+1)/\beta$ nonintegral, which gives the stated range. The unique integral solution in the block is N itself, whence the additional 1; passing through $m = 2^x$ turns the exponent block $[N, N+1)$ into the magnitude block $[2^N, 2^{N+1})$ bijectively. $\square$

Status: [\[kernel-verified Lean\]](#)

`ExactCounting` proves the explicit ceiling-image parametrization, `twoBaseNonintegerSolutionsInUnitBlock_ncard`, `twoBaseNaturalCandidatesInDyadicBlock_card`, and the exact finite-sum identity `cumulativeTwoBaseCandidateCount_eq`. `QuadraticAsymptotic` identifies that cumulative count with the candidate count below 2^T and proves the sharp bounds and the limit in (5).

In particular a failed conjecture would produce only linearly many candidates in the logarithmic block index,

$$\#(\mathcal{C} \cap [X, 2X]) \asymp \log X,$$

and the kernel-verified zero-free region gives $\beta \geq 12$, hence $\beta > 12$ because β is irrational. The number of *nonintegral* solutions in the block is therefore at most $\lfloor N/12 \rfloor$, while the full solution count — equivalently the dyadic candidate count — is at most $1 + \lfloor N/12 \rfloor$. Under the directed-rounding scan these bounds improve respectively to $\lfloor N/27 \rfloor$ and $1 + \lfloor N/27 \rfloor$, at the weaker [\[software computation\]](#) trust level. Summing over blocks, the cumulative candidate count below 2^T is

$$C(T) = T + \sum_{r=1}^T \left\lfloor \frac{r}{\beta} \right\rfloor = \frac{T^2}{2\beta} + O(T). \quad (5)$$

More precisely, Lean proves the two-sided bound and the limit

$$\frac{T(T+1)}{2\beta} \leq C(T) \leq \frac{T(T+1)}{2\beta} + T, \quad \frac{C(T)}{T^2} \longrightarrow \frac{1}{2\beta}.$$

This is conditionally sharper than the unconditional verified bounds $((\log_2 N)^2$ and $\log_2 N \cdot \log_3 N$ of Section 4) and matches them in leading order: with $\beta > 12$ the leading coefficient $1/(2\beta)$ is below $1/24$.

There is also an unconditional growth reformulation, obtained by comparing the two cases. Under the conjecture the count below 2^T is exactly T ; under failure the preceding limit is the strictly positive constant $1/(2\beta)$. Hence

$$\text{Conjecture 2.1 holds} \iff \frac{\#(\mathcal{C} \cap [1, 2^T])}{T^2} \longrightarrow 0, \quad (6)$$

which is `alaogluErdosConjecture_iff_candidate_count_below_two_pow_subquadratic`. [\[kernel-verified Lean\]](#)

This exact population is the correct benchmark for every proposed density or determinant argument. A method that proves only $O((\log X)^2)$ candidates in $[X, 2X]$, for instance, is a valid unconditional theorem but cannot contradict Theorem 5.8: it is weaker than the conditional truth it would need to refute.

5.6 Blockwise irrational rotation and equidistribution

By Theorem 5.8 the fractional part of $x_{N,k}$ is $\{k\beta\}$, so a high unit block contains exactly an initial segment

$$\{\beta\}, \{2\beta\}, \dots, \{K_N\beta\}, \quad K_N = \left\lfloor \frac{N+1}{\beta} \right\rfloor,$$

of one irrational rotation orbit, and $K_N \rightarrow \infty$.

Theorem 5.9 (Finite Fourier tests on exact blocks). *Assume a counterexample exists and let β be its primitive generator. For every finitely supported coefficient function $c : \mathbb{Z} \rightarrow \mathbb{C}$, put*

$$P_c(t) = \sum_{h \in \mathbb{Z}} c(h) e^{2\pi i h t}.$$

Then, on the exact ceiling-parametrized points $x_{N,k} = \lceil N - k\beta \rceil + k\beta$,

$$\frac{1}{K_N} \sum_{k=1}^{K_N} P_c(x_{N,k}) \longrightarrow c(0) \quad (N \rightarrow \infty).$$

Status: [\[kernel-verified Lean\]](#)

`BlockWeyl` proves vanishing of every nonzero normalized Fourier mode on the exact block points, and `BlockEquidistribution` proves `tendsto_primitiveUnitBlockTrigonometricAverage` for every finitely supported coefficient function, with limit $c = 0$. This is convergence against finite trigonometric polynomials; it is not yet a theorem for arbitrary continuous tests.

Theorem 5.10 (Blockwise Weyl equidistribution). *Assume a counterexample exists. For every interval $I \subseteq [0, 1]$ whose boundary has measure zero,*

$$\frac{\#\{x \in (\mathcal{S} \setminus \mathbb{N}_0) \cap [N, N+1) : \{x\} \in I\}}{\#((\mathcal{S} \setminus \mathbb{N}_0) \cap [N, N+1))} \longrightarrow |I| \quad (N \rightarrow \infty).$$

Proof. By Theorem 5.8 the relevant fractional parts are $\{\beta\}, \{2\beta\}, \dots, \{K_N\beta\}$ with $K_N = \lfloor (N+1)/\beta \rfloor \rightarrow \infty$. For every nonzero integer h the Weyl sum is a finite geometric series,

$$\frac{1}{K} \sum_{k=1}^K e^{2\pi i h k \beta} = \frac{e^{2\pi i h \beta} (1 - e^{2\pi i h K \beta})}{K(1 - e^{2\pi i h \beta})} \longrightarrow 0,$$

since β is irrational and hence $e^{2\pi i h \beta} \neq 1$. Weyl’s criterion concludes. $\square$

Status: [\[paper proof in this report\]](#)

The finite-trigonometric-polynomial input is kernel-verified by Theorem 5.9, but the Stone–Weierstrass and Weyl-criterion extension to all continuous tests, and then to interval indicators, has not been formalized. No effective discrepancy rate follows without further Diophantine input on β : the geometric-sum proof controls each fixed Fourier mode but not uniformly over many modes.

This regularity is attractive, and it is worth stating plainly why it does not by itself decide anything. The integrality condition being tested is not a fixed positive-measure target in the circle; it is an exponentially shrinking target, since the relevant approximation quality must improve like the reciprocal of the integer sizes involved. Equidistribution at a fixed scale is simply the wrong instrument for a target whose measure tends to zero geometrically, and that mismatch recurs in the analytic approaches examined later in this report.

6 Equivalent one-number and radical formulations

The exact conditional structure of Section 5 compresses a statement quantified over all real x into a statement about a single real number at a time. This section records that compression in its two kernel-verified forms — a localization at the prime 3 and a radical-algebraicity normalization — and then names the exact arithmetic obstruction that remains.

6.1 Localization at the prime three

Theorem 6.1 (Localization reformulation). *Conjecture 2.1 is false if and only if there exists an odd integer $u > 1$ with*

$$u^\vartheta \in \mathbb{Z}[1/3],$$

equivalently, an odd $u > 1$ and $a \in \mathbb{N}_0$ with $3^a u^\vartheta \in \mathbb{N}$.

Proof. If x is a counterexample, write $2^x = 2^a u$ with u odd; since 2^x is not a power of two, $u > 1$. Using $(2^a)^\vartheta = 3^a$,

$$u^\vartheta = \frac{(2^x)^\vartheta}{(2^a)^\vartheta} = \frac{3^x}{3^a} \in \mathbb{Z}[1/3].$$

Conversely, if $u^\vartheta = B/3^a$ with $B \in \mathbb{N}$, set $x = a + \log_2 u$; then $2^x = 2^a u \in \mathbb{N}$ and $3^x = 3^a u^\vartheta = B \in \mathbb{N}$, and $x \notin \mathbb{Z}$ since otherwise the odd $u > 1$ would be a power of two. $\square$

Status: [\[kernel-verified Lean\]](#) —

`alaogluErdosConjecture_iff_odd_rpow_not_localized` in `Localization` and `alaogluErdosConjecture_iff_oddLocalizedRadicalExclusion` in `RadicalReformulation`; the candidate-level variant `alaogluErdosConjecture_iff_no_odd_multiple_candidate`, also in `Localization`, is proved independently and phrases the same content in terms of admissible candidate bases. The equivalence is kernel-verified; the localized-radical exclusion on its right-hand side remains the open conjecture. The allowance of 3-power denominators is exactly the freedom created by shifting the exponent by an integer.

The denominator restriction to powers of 3 is not cosmetic. Only the prime 3 can occur in the denominator, because the exponent shift $x \mapsto x + 1$ multiplies 2^x by 2 and 3^x by 3; the odd part u absorbs the 2-adic freedom completely and $\mathbb{Z}[1/3]$ absorbs the remaining 3-adic freedom. Everything else has been eliminated.

6.2 Reduction to power-primitive odd bases

An odd $u > 1$ that is a proper power, $u = w^d$ with w odd and power-primitive, satisfies $u^\vartheta = (w^\vartheta)^d$. Unique primitive-power decomposition therefore lets the quantifier in Theorem 6.1 range over power-primitive bases only, at the cost of an extra exponent $d \geq 1$.

Status: [\[kernel-verified Lean\]](#) — `PrimitiveRadical` proves unique primitive-power decomposition of every integer $u > 1$ together with the exact equivalence `alaogluErdosConjecture_iff_primitiveOddLocalizedRadicalExclusion`. Proper-power choices of u introduce no genuinely new cases.

6.3 Canonical radical algebraicity

In the notation of the least generator of Section 5, a counterexample produces natural numbers w, d, a, A with $w > 1$ odd and power-primitive, $d > 0$, and

$$\beta = a + d \log_2 w, \quad 2^\beta = 2^a w^d, \quad 3^\beta = A,$$

normalized so that $a = 0$ or $3 \nmid A$. Setting

$$E := w^\vartheta = 3^{\log_2 w},$$

one has $E^d = A/3^a \in \mathbb{Q}$, and d is the least positive exponent for which E^d is rational. Minimality of d has an exact algebraic consequence.

Proposition 6.2 (Exact radical degree). *The binomial $X^d - A/3^a$ is irreducible over $\mathbb{Q}$, it is the minimal polynomial of E , and E is algebraic of exact degree d .*

Proof. For a positive rational q , the binomial $X^d - q$ can be reducible only if q is a p -th power in $\mathbb{Q}$ for some prime $p \mid d$ (the additional $4 \mid d$ exceptional case concerns -4 times a fourth power and is irrelevant for $q > 0$). If $q = r^p$ with $p \mid d$, the positive real $E^{d/p}$ is the positive p -th root of q , hence rational, contradicting minimality of d . $\square$

Status: [\[kernel-verified Lean\]](#) —

`irreducible_X_pow_sub_C_of_least_rational_power` in `RadicalDegree` proves the general norm argument, and `oddCoreRpow_exact_radical_degree` specializes it to the least normalized radical index. It identifies the minimal polynomial with $X^d - A/3^a$ and proves that the degree is exactly d .

A counterexample therefore produces the very special algebraic number $E = 3^{\log_2 w}$, reached from a transcendental exponent. Taking logarithms of $E^d = A/3^a$:

$$d \log w \log 3 = (\log A - a \log 3) \log 2. \quad (7)$$

This is a *bilinear* relation among logarithms of integers: each side is a product of two logarithms, not a linear form in logarithms with algebraic coefficients. Baker’s theorem [3] bounds linear forms $\sum \alpha_i \log \gamma_i$ with algebraic α_i away from zero, and no substitution turns (7) into such a form, because the coefficient of $\log w$ on the left is $d \log 3$, itself a logarithm rather than an algebraic number. The four exponentials conjecture does forbid the relation. Equation (7) is thus a precise description of the transcendence-theoretic wall, not a disguised elementary factorization problem.

6.4 Where the wall stands in the smallest case

The smallest structural question already meets the open core. A 3-smooth candidate $m = 2^a 3^b$ with $b \geq 1$ would require $3^a E_3^b \in \mathbb{N}$ for

$$E_3 = 3^{\log_2 3} = e^{(\log 3)^2 / \log 2} = 5.7045224946911176 \dots,$$

and even the irrationality of E_3 is open: it is a “three-logarithms” quantity outside the reach of Gelfond–Schneider, Baker’s method, and the six exponentials theorem. No currently known technique excludes the odd core $w = 3$; this is why the corpus emphasizes counting, structure, and verified finite regions rather than infinite families of exclusions, and why Conjecture 6.3 below is the honest name of the obstruction.

Conjecture 6.3 (Radical four-exponentials target). *For every odd integer $u > 1$, $u^{\log_2 3} \notin \mathbb{Z}[1/3]$. Equivalently, for every power-primitive odd $w > 1$ and every $d \geq 1$, $(3^{\log_2 w})^d \notin \mathbb{Z}[1/3]$.*

By Theorem 6.1 and the primitive reduction above, both formulations are exactly equivalent to Conjecture 2.1; neither is proved. A proof of the special case $u = 3$ — even of the weaker assertion that $3^{\log_2 3}$ is irrational — by any genuinely new mechanism would already be a meaningful breakthrough.

7 The exact four-exponentials barrier

7.1 The two-by-two exponential matrix

Suppose x is a nonintegral solution. Consider the row vector of parameters $(1, x)$ and the column vector of parameters $(\log 2, \log 3)$. The four exponentials formed from the products are

$$\exp(1 \cdot \log 2) = 2, \quad \exp(1 \cdot \log 3) = 3, \quad \exp(x \log 2) = 2^x, \quad \exp(x \log 3) = 3^x,$$

all four algebraic — indeed positive integers. The two row parameters $1, x$ are $\mathbb{Q}$ -linearly independent because x is irrational: a rational nonintegral exponent is excluded by the rational-exponent argument of Section 2. The two column parameters $\log 2, \log 3$ are $\mathbb{Q}$ -linearly independent because $2^a 3^b = 1$ forces $a = b = 0$, that is, no nonzero powers of 2 and 3 agree. The real four exponentials conjecture predicts that at least one of the four displayed values must be transcendental, a contradiction [10].

Status: [\[kernel-verified Lean\]](#) —

`alaogluErdosConjecture_of_realFourExponentialsConjecture` in

`FourExponentialsReduction` formalizes exactly this implication. It does not assume four exponentials as an axiom: the conjectural statement is defined as a proposition and the implication is proved conditionally on it. Four exponentials itself is open, so this is context and diagnosis, not a proof of Conjecture 2.1.

The reduction is a diagnostic and not merely background. Any proposed elementary argument must exploit integrality, positivity, or arithmetic structure beyond bare algebraicity of those four values, or else it would prove a substantial case of the four exponentials conjecture by another name.

7.2 Why six exponentials does not specialize

The six exponentials theorem applies to two independent row parameters together with three independent column parameters, or to three rows and two columns. Here only two canonical logarithmic directions are available. Adding a third base 5 supplies the missing direction and yields the kernel-verified three-base theorem

$$2^x, 3^x, 5^x \in \mathbb{Z} \implies x \in \mathbb{Z},$$

proved by an interpolation determinant. In the two-base problem, any additional base built from powers of 2 and 3 has logarithm in the rational span of $\log 2$ and $\log 3$ and therefore does not create an independent column.

The dimension count is what decides the outcome. With three independent row directions and two column directions the analytic size estimate beats the arithmetic lower bound $|\det| \geq 1$; in the 2×2 configuration the same balance lands exactly on the classical boundary, where the required decay below one fails. A successful two-base determinant argument would, by construction, be a new four-exponentials argument.

The classification of auxiliary bases makes the obstruction precise: `ThreeSmooth` and the denominator-controlled results in `RationalThirdBase` show that an auxiliary base with a prime factor outside $\{2, 3\}$ would supply a genuinely new column, but its corresponding exponential value is not forced to be an integer — or even algebraic — by the original hypotheses. Bases supported only on 2 and 3 do not increase the transcendence rank.

7.3 Integrality is stronger than algebraicity, but not by enough

Four exponentials asks only for algebraicity. The values here are integers, so there is extra discrete structure available, and the determinant machinery of Section 19 and Section 20 extracts some of it: the relevant determinants have integer entries, so a nonzero determinant is a nonzero integer and therefore satisfies $|\det| \geq 1$.

That inequality is the whole of the extra archimedean information. A nonzero integer has absolute value at least one, and nothing in the integrality hypotheses improves the constant 1 by itself. To cross the four-exponentials threshold one would need a strictly stronger arithmetic input, for instance:

- a proof that the determinants are divisible by a large integer, replacing $|\det| \geq 1$ by $|\det| \geq D$ with D growing in the size parameters; or
- a proof that a suitable common denominator is anomalously small, improving the arithmetic side rather than the analytic side.

Both routes ask for divisibility rather than size, which is precisely why the p -adic direction is given priority in the research program of Section 27.

Status: [\[paper proof in this report\]](#) for the unconditional paper proofs of the determinant hierarchy. [\[kernel-verified Lean\]](#) for the accepted algebraic, integrality, nonvanishing, and positivity layers in `ArbitraryNodeDeterminant`, `SemigroupDeterminant`, `MixedExponentSemigroup`, and `OrderedChamberPositivity`. The quantitative Lean endpoints carry their still-missing calculus inputs as explicit theorem hypotheses rather than axioms. The reduction of Conjecture 2.1 to four exponentials is [\[kernel-verified Lean\]](#), but four exponentials is open and the conjecture remains open with it.

8 Rational-function rigidity and closure properties

The exact primitive-generator theorem of Section 5 determines the additive structure of $\mathcal{S}$ completely under failure of Conjecture 2.1. It has a consequence that deserves to be isolated: a hypothetical solution monoid is not merely additively free of rank two, it is *rigid* against every nonlinear rational operation over the algebraic numbers. Addition and multiplication by naturals are the only operations that stay inside $\mathcal{S}$, and every apparent exception collapses to an identity in one indeterminate. The results below are exact equivalences, several of them unconditional, and they translate directly into closure reformulations of the conjecture itself.

8.1 Evaluation injectivity at the generator

Throughout this section, β denotes the canonical least nonintegral solution supplied by the generator theorem of Section 5, so that

$$\mathcal{S} = \{n + k\beta : n, k \in \mathbb{N}_0\}$$

with unique coordinates n, k . The transcendence dichotomy of Section 4 makes β transcendental over $\mathbb{Q}$. The rational-function statements need the slightly stronger form over the algebraic closure.

Lemma 8.1 (Evaluation is injective over $\overline{\mathbb{Q}}$). *Let $\gamma \in \mathbb{C}$ be transcendental over $\mathbb{Q}$. Then no nonzero $P \in \overline{\mathbb{Q}}[T]$ satisfies $P(\gamma) = 0$; equivalently, the evaluation homomorphism $\overline{\mathbb{Q}}[T] \rightarrow \mathbb{C}$, $P \mapsto P(\gamma)$, is injective.*

Proof. If $P(\gamma) = 0$ for some nonzero $P \in \overline{\mathbb{Q}}[T]$, then γ is algebraic over $\overline{\mathbb{Q}}$. Since $\overline{\mathbb{Q}}/\mathbb{Q}$ is an algebraic extension, the tower property makes γ algebraic over $\mathbb{Q}$, contradicting the hypothesis. Injectivity follows by applying this to differences. $\square$

Status: [\[paper proof in this report\]](#) — the transcendence of β itself is [\[kernel-verified Lean\]](#) (Section 4, Transcendence); the only new content here is the standard tower argument that upgrades transcendence over $\mathbb{Q}$ to transcendence over $\overline{\mathbb{Q}}$.

8.2 The rational-function intersection theorem

Theorem 8.2 (Rational-function rigidity at the primitive generator). *Assume Conjecture 2.1 is false and let β be the canonical generator. Let $R(T) \in \overline{\mathbb{Q}}(T)$ have no pole at β . Then*

$$R(\beta) \in \mathcal{S} \iff R(T) = n + kT \text{ in } \overline{\mathbb{Q}}(T) \text{ for some } n, k \in \mathbb{N}_0.$$

Proof. The reverse implication is additive closure of $\mathcal{S}$ together with closure under multiplication by naturals: $R(\beta) = n + k\beta \in \mathcal{S}$.

For the forward implication write $R = P/Q$ with $P, Q \in \overline{\mathbb{Q}}[T]$ and $Q(\beta) \neq 0$. If $R(\beta) \in \mathcal{S}$, the exact coordinates give $R(\beta) = n + k\beta$ for unique $n, k \in \mathbb{N}_0$, hence

$$P(\beta) - (n + k\beta)Q(\beta) = 0.$$

The left-hand side is the value at β of the polynomial $P(T) - (n + kT)Q(T) \in \overline{\mathbb{Q}}[T]$. By Lemma 8.1 that polynomial is identically zero, so $P(T) = (n + kT)Q(T)$ and $R(T) = n + kT$ in $\overline{\mathbb{Q}}(T)$. $\square$

Status: [\[paper proof in this report\]](#)

The theorem allows poles away from β , permits arbitrary cancellation, and takes the full algebraic closure as coefficient field. It can be read as the exact intersection

$$\mathcal{S} \cap \overline{\mathbb{Q}}(\beta) = \mathbb{N}_0 + \mathbb{N}_0\beta, \tag{8}$$

with the extra information that every rational expression realizing an element of that intersection is *identically* the corresponding affine expression — there is no accidental representation.

Corollary 8.3 (Rational self-maps of the solution set). *Under failure, the rational functions over $\overline{\mathbb{Q}}$ that are regular on $\mathcal{S}$ and map $\mathcal{S}$ into itself are exactly the maps*

$$T \mapsto n + kT, \quad n, k \in \mathbb{N}_0.$$

Proof. Such a function in particular sends β into $\mathcal{S}$, so Theorem 8.2 identifies it as $n + kT$ with $n, k \in \mathbb{N}_0$. Conversely, translation by a natural and dilation by a natural preserve $\mathbb{N}_0 + \mathbb{N}_0\beta$. $\square$

Status: [paper proof in this report]

Proposition 8.4 (Multivariate reduction to one indeterminate). *Assume failure, let $x_i = n_i + k_i\beta \in \mathcal{S}$ for $i = 1, \dots, r$, and let $F \in \overline{\mathbb{Q}}(X_1, \dots, X_r)$ be regular at $(x_1, \dots, x_r)$. Then*

$$F(x_1, \dots, x_r) \in \mathcal{S}$$

if and only if the one-variable specialization

$$G(T) := F(n_1 + k_1T, \dots, n_r + k_rT)$$

is identically $p + qT$ in $\overline{\mathbb{Q}}(T)$ for some $p, q \in \mathbb{N}_0$.

Proof. Write $F = P/Q$ with $P, Q \in \overline{\mathbb{Q}}[X_1, \dots, X_r]$ and $Q(x_1, \dots, x_r) \neq 0$. The specialized denominator $Q(n_1 + k_1T, \dots, n_r + k_rT)$ is a polynomial in T that does not vanish at $T = \beta$, hence is not the zero polynomial. Therefore G is a well-defined element of $\overline{\mathbb{Q}}(T)$ with no pole at β , and $G(\beta) = F(x_1, \dots, x_r)$. Now apply Theorem 8.2. $\square$

Status: [paper proof in this report]

Every rational operation on finitely many hypothetical solutions is thus reduced to an elementary identity between polynomials in a single indeterminate, with coordinates in $\mathbb{N}_0$. No analytic input is required beyond the transcendence already verified by the kernel.

8.3 Exact multiplication criterion

The next statement needs no assumption at all: the branch in which the conjecture holds is trivial, and the branch in which it fails is governed by the generator.

Theorem 8.5 (Products of two solutions). *For all $x, y \in \mathcal{S}$, unconditionally,*

$$xy \in \mathcal{S} \iff x \in \mathbb{N}_0 \text{ or } y \in \mathbb{N}_0.$$

Proof. If one factor, say x , lies in $\mathbb{N}_0$, then xy is a natural multiple of a solution and therefore a solution; this direction is unconditional.

For the converse, split on Conjecture 2.1. If the conjecture holds then $\mathcal{S} = \mathbb{N}_0$ and the right-hand side is automatically true, so there is nothing to prove. Under failure, write the unique coordinates

$$x = n + k\beta, \quad y = m + \ell\beta, \quad n, k, m, \ell \in \mathbb{N}_0.$$

If neither factor is integral then $k, \ell > 0$ and

$$xy = nm + (n\ell + mk)\beta + k\ell\beta^2,$$

with $k\ell > 0$. Membership $xy \in \mathcal{S}$ would force $xy = r + s\beta$ for some $r, s \in \mathbb{N}_0$, producing a nonzero polynomial of degree exactly two over $\mathbb{Q}$ vanishing at the transcendental number β . Equivalently, apply Theorem 8.2 to $R(T) = (n + kT)(m + \ell T)$, which is not affine. Both routes are impossible. $\square$

Status: [\[paper proof in this report\]](#) — note that only transcendence over $\mathbb{Q}$ is used, which is [\[kernel-verified Lean\]](#); the appeal to $\mathbb{Q}$ in Theorem 8.2 is a convenience, not a necessity.

Corollary 8.6 (Finite products). *Let $x_1, \dots, x_r \in \mathcal{S}$ be positive. Then*

$$\prod_{i=1}^r x_i \in \mathcal{S} \iff \text{at most one } x_i \text{ is nonintegral.}$$

If a zero factor is allowed, the product is $0 \in \mathcal{S}$ and the criterion is vacuous.

Proof. If the conjecture holds, every x_i is integral and both sides are true. Under failure, the positive integral factors multiply to a positive natural, so after removing them the product is a natural multiple of a product of nonintegral solutions. Writing each remaining factor as $n_i + k_i\beta$ with $k_i > 0$ and expanding, the result is a polynomial in β whose degree equals the number of nonintegral factors and whose leading coefficient is positive. Degree at least two is excluded by Theorem 8.2, while degree zero or one is preserved by natural dilation and translation. $\square$

Status: [\[paper proof in this report\]](#)

Corollary 8.7 (Internal power test). *For every $x \in \mathcal{S}$ and every integer $r \geq 2$,*

$$x^r \in \mathcal{S} \iff x \in \mathbb{N}_0.$$

In particular $x \in \mathbb{N}_0$ holds exactly when $x^2 \in \mathcal{S}$: a single squaring test detects integrality.

Proof. Apply Theorem 8.5 with $y = x$ and induct, or apply Corollary 8.6 to r equal factors, noting that $x = 0$ satisfies both sides. $\square$

Status: [\[paper proof in this report\]](#)

8.4 Exact quotient criterion

The same evaluation argument classifies division. Division is more delicate than multiplication because even the true solution set $\mathbb{N}_0$ is not closed under arbitrary quotients, so the criterion cannot be a clean dichotomy; it is instead an exact divisibility condition on coordinates.

Theorem 8.8 (Quotients of solutions). *Assume Conjecture 2.1 is false and write the unique coordinates*

$$x = n + k\beta, \quad y = m + \ell\beta > 0.$$

Exactly one of the two cases $\ell = 0$ and $\ell > 0$ applies, and in that case $x/y \in \mathcal{S}$ if and only if the corresponding condition holds:

- (a) $\ell = 0$, so that $y = m$ is a positive natural, and $m \mid n$ and $m \mid k$;
- (b) $\ell > 0$, and $x = qy$ for some $q \in \mathbb{N}_0$, equivalently $n = qm$ and $k = q\ell$.

In particular, if x and y are both nonintegral, then

$$\frac{x}{y} \in \mathcal{S} \iff \frac{x}{y} \in \mathbb{N}_0.$$

Proof. Suppose $x/y = r + s\beta$ with $r, s \in \mathbb{N}_0$. Clearing the denominator gives

$$n + k\beta = mr + (ms + \ell r)\beta + \ell s \beta^2.$$

The difference of the two sides is a polynomial in β with rational coefficients, so transcendence of β forces every coefficient to vanish:

$$\ell s = 0, \quad k = ms + \ell r, \quad n = mr.$$

If $\ell = 0$ then $y = m$ and $y > 0$ gives $m \geq 1$; the equations read $n = mr$ and $k = ms$, which is precisely $m \mid n$ and $m \mid k$. If $\ell > 0$ then $s = 0$, so $n = mr$ and $k = \ell r$, that is $x = ry$ with $r \in \mathbb{N}_0$. The converses follow by direct substitution.

For the final claim, both x and y nonintegral means $k, \ell > 0$, so case (b) applies and $x/y = r \in \mathbb{N}_0$. $\square$

Status: [\[paper proof in this report\]](#)

Corollary 8.9 (Reciprocals). *Unconditionally, for $x \in \mathcal{S}$ with $x > 0$,*

$$\frac{1}{x} \in \mathcal{S} \iff x = 1.$$

Proof. If the conjecture holds then $\mathcal{S} = \mathbb{N}_0$ and the claim is the elementary fact that a positive natural with natural reciprocal equals 1. Under failure apply Theorem 8.8 with numerator $1 = 1 + 0 \cdot \beta$ and denominator $y = x$. If $x = m$ is integral, case (a) gives $m \mid 1$, so $m = 1$. If x is nonintegral, case (b) gives $1 = qx$ for some $q \in \mathbb{N}_0$; here $q = 0$ is impossible, and $q \geq 1$ would make $x = 1/q$ rational, contradicting the irrationality of every nonintegral solution. $\square$

Status: [\[paper proof in this report\]](#)

A hypothetical solution monoid is therefore additively rich — a free additive monoid of rank two — and almost maximally hostile to every multiplicative operation: products, powers, reciprocals and quotients all escape $\mathcal{S}$ except in the degenerate cases forced by the additive structure itself.

8.5 Closure formulations of the conjecture

Because the multiplication criterion is unconditional, it converts directly into equivalent statements of the conjecture in terms of closure.

Theorem 8.10 (Multiplicative closure equivalences). *The following are equivalent.*

- (i) *Conjecture 2.1 is true.*
- (ii) *$\mathcal{S}$ is closed under multiplication.*

(iii) $\mathcal{S}$ is closed under squaring.

(iv) For every real x ,

$$2^x, 3^x \in \mathbb{Z} \implies 2^{x^2}, 3^{x^2} \in \mathbb{Z}.$$

Proof. If the conjecture holds then $\mathcal{S} = \mathbb{N}_0$, which is closed under multiplication, so (i) implies (ii); and (ii) trivially implies (iii). Statements (iii) and (iv) are the same assertion, written once for $\mathcal{S}$ and once for the defining integrality conditions. Finally, assume (iii). For any $x \in \mathcal{S}$ we get $x^2 \in \mathcal{S}$, so Corollary 8.7 forces $x \in \mathbb{N}_0$. Hence $\mathcal{S} = \mathbb{N}_0$, which is Conjecture 2.1. $\square$

Status: [paper proof in this report]

The candidate formulation of Section 2 carries the same content in multiplicative language. For positive reals define the commutative operation

$$m \star n := 2^{(\log_2 m)(\log_2 n)} = m^{\log_2 n} = n^{\log_2 m}. \quad (9)$$

If $m = 2^x$ and $n = 2^y$, then $m \star n = 2^{xy}$, so $\star$ is exactly the transport of multiplication of exponents through the exponent–candidate bijection of Section 2.

Corollary 8.11 (Candidate star-product). *For $m, n \in \mathcal{C}$,*

$$m \star n \in \mathcal{C} \iff m \text{ is a power of 2 or } n \text{ is a power of 2}.$$

Consequently Conjecture 2.1 is equivalent both to closure of $\mathcal{C}$ under $\star$ and to the apparently weaker diagonal condition

$$m \in \mathcal{C} \implies m \star m \in \mathcal{C}.$$

Proof. Put $x = \log_2 m$ and $y = \log_2 n$, so that $x, y \in \mathcal{S}$ and $m \star n = 2^{xy}$. Membership $m \star n \in \mathcal{C}$ is equivalent to $xy \in \mathcal{S}$, and m is a power of two exactly when $x \in \mathbb{N}_0$; now apply Theorem 8.5. The two closure equivalences are the image of Theorem 8.10(ii) and (iii) under the same bijection. $\square$

Status: [paper proof in this report]

These are genuine equivalences and not a proof. There is at present no independent mechanism forcing nonlinear closure: nothing in the archimedean, valuation-theoretic, or determinant toolkits of Sections 4, 19 and 20 supplies a reason for $\mathcal{C}$ to be closed under $\star$. Their value is diagnostic — they identify the missing property as a single crisp failure of semiring-like structure, rather than as a quantitative deficit.

8.6 Formalization cost

The multiplication, power and quotient criteria are inexpensive Lean targets, and a route exists that avoids building a general rational-function library. For a new module — `RationalRigidity` would be a natural name — the plan is:

1. case split on `AlaogluErdosConjecture`, discharging the true branch by $\mathcal{S} = \mathbb{N}_0$;

2. in the failure branch obtain the canonical generator and the unique coordinates from the packaged statement

`exists_canonical_primitiveGenerator_of_not_alaogluErdosConjecture`

in the module `PrimitiveGenerator`;

3. expand the relevant equality as a polynomial of degree at most two in β with rational coefficients;
4. invoke the kernel-verified transcendence in `Transcendence` to force every coefficient to vanish;
5. discharge the remaining natural-number divisibility and coordinate conditions with `omega`.

Only the full rational-function theorem, Theorem 8.2, needs evaluation injectivity over $\overline{\mathbb{Q}}$ and hence Lemma 8.1; the multiplication theorem, the power test, the quotient criterion and the closure equivalences need nothing beyond transcendence over $\mathbb{Q}$, which is already accepted by the kernel.

Status: Planned; no draft yet. The mathematics is [\[paper proof in this report\]](#) throughout this section, and the prerequisites are [\[kernel-verified Lean\]](#).

8.7 What the rigidity statements buy

The practical payoff is that they sharply restrict how a counterexample could be manufactured from a smaller one. A natural strategy for producing new solutions from old ones is to apply some algebraic operation — square a solution, multiply two of them, take a ratio, invert, or substitute solutions into a rational function — and hope to land back in $\mathcal{S}$, either to run a descent or to build an infinite family with better arithmetic properties than the canonical generator. Equation (8) closes that door completely: the only rational functions over $\overline{\mathbb{Q}}$ that map solutions to solutions are the affine maps $T \mapsto n + kT$ with natural coefficients, which reproduce nothing that the additive monoid $\mathbb{N}_0 + \mathbb{N}_0\beta$ did not already contain. Any construction of a new counterexample from an old one must therefore be transcendental in an essential way, or must leave the solution set and return by a route that is not rational over the algebraic numbers. In the other direction, the closure equivalences of Theorem 8.10 and Corollary 8.11 convert the conjecture into a purely structural question — is $\mathcal{C}$ closed under $\star$? — which is attractive precisely because it has no quantitative parameter to optimize, and which is a warning that no counting or equidistribution bound alone will produce the answer. Finally, all of these statements are short, elementary consequences of results already accepted by the kernel, which makes them the cheapest genuine additions to the verified corpus that this report proposes.

9 Three-smooth outputs and transcendental product exponents

Call a positive integer *three-smooth* when it has no prime factor outside $\{2, 3\}$, that is, when it equals $2^u 3^v$ for some $u, v \in \mathbb{N}_0$. The kernel-verified classification of auxiliary bases recorded in Section 4 settles which bases B satisfy $B^x \in \mathbb{Z}$ for a *fixed* nonintegral solution x : exactly the three-smooth ones. This section turns that classification onto the outputs $2^x, 3^x$ themselves, and then onto second-level exponents xy formed from two solutions.

Two distinct things come out, and they have different provenance. The first is an unconditional integrality test: a solution is integral precisely when its own two outputs are three-smooth. It uses only the kernel-verified transcendence of ϑ and unique factorization, so it is cheap to add to the Lean development. The second is an upgrade of “not an integer” to “transcendental”, obtained from the six exponentials theorem. That upgrade converts several apparently weaker *algebraicity* statements into exact reformulations of Conjecture 2.1, at the cost of a classical input that the repository has not formalized in general form.

9.1 The two outputs cannot both be three-smooth

Theorem 9.1 (Three-smooth output test). *For every $x \in \mathcal{S}$ the following are equivalent:*

- (i) $x \in \mathbb{N}_0$;
- (ii) both 2^x and 3^x are three-smooth;
- (iii) 6^x is three-smooth.

Proof. Condition (i) gives (ii) immediately, since 2^x and 3^x are then powers of 2 and 3 respectively. Conditions (ii) and (iii) are equivalent because 2^x and 3^x are positive integers with product 6^x , and a product of positive integers is three-smooth exactly when both factors are.

For (ii) $\Rightarrow$ (i), write

$$2^x = 2^p 3^q, \quad 3^x = 2^r 3^s$$

with $p, q, r, s \in \mathbb{N}_0$. Taking logarithms and dividing by $\log 2$ and $\log 3$ respectively gives

$$x = p + q\vartheta, \quad x = s + \frac{r}{\vartheta},$$

where $\vartheta = \log_2 3$. Eliminating x and clearing the denominator yields the quadratic relation

$$q\vartheta^2 + (p - s)\vartheta - r = 0 \tag{10}$$

with rational — indeed integer — coefficients. The repository proves ϑ transcendental, so every coefficient in (10) vanishes: $q = 0$, $p = s$, and $r = 0$. Hence $2^x = 2^p$ and $x = p \in \mathbb{N}_0$. $\square$

Status: [\[paper proof in this report\]](#)

The statement is unconditional and does not use the conditional generator, the exact solution monoid, or any six-exponentials input. Its only nonelementary ingredient is the kernel-verified transcendence of ϑ ,

`transcendental_logThreeDivLogTwo`

in **Transcendence**, together with unique factorization. A Lean proof is expected to be short; see Section 28 and the skeleton at the end of this section.

The test is sharp in a useful direction: it rules out one of the four a priori smoothness patterns of the least output pair.

Corollary 9.2 (Primitive-core trichotomy). *Assume Conjecture 2.1 fails, and use the simultaneous least-output normalization of Section 5,*

$$M = 2^\beta = 2^a w^d, \quad A = 3^\beta = 3^b v^e,$$

with $w > 1$ odd and power-primitive, $v > 1$ power-primitive and coprime to 3, and $d, e > 0$. Then at most one of $w = 3$ and $v = 2$ holds, and exactly one of the following three patterns occurs:

M	A	primitive cores
three-smooth	not three-smooth	$w = 3, v \neq 2$
not three-smooth	three-smooth	$w \neq 3, v = 2$
not three-smooth	not three-smooth	$w \neq 3, v \neq 2$

The missing fourth pattern, both outputs three-smooth, is impossible.

Proof. An odd, three-smooth, power-primitive integer greater than one is a power of 3, and primitivity forces the exponent to be 1; hence M is three-smooth exactly when $w = 3$. Symmetrically a three-smooth power-primitive integer greater than one that is coprime to 3 is exactly 2, so A is three-smooth exactly when $v = 2$. The least nonintegral solution β lies in $\mathcal{S} \setminus \mathbb{N}_0$, so Theorem 9.1 forbids both outputs from being three-smooth, which is the exclusion $\neg(w = 3 \wedge v = 2)$. $\square$

Status: [paper proof in this report]

The normalization itself, including power-primitivity of w and v , the depth dichotomy $a = 0$ or $b = 0$, and the gcd-one condition, is [kernel-verified Lean] in `OutputNormalization` (Section 5). Only the exclusion of the fourth pattern is new here.

The two exceptional rows are precisely the cases the corpus already identifies as hardest: $w = 3$ makes 3^β the sharp radical target of Section 6, and $v = 2$ is its base-swapped mirror with $2^{1/\beta} = 2^{\log_3 2}$. The third row is generic and carries two independent outside-prime supports.

9.2 Third-base transcendence from six exponentials

The remaining statements of this section rest on one published transcendence theorem, which we state in the form we use.

Theorem 9.3 (Six exponentials theorem; Siegel, Lang, Ramachandra). *Let u_1, u_2 be complex numbers linearly independent over $\mathbb{Q}$, and let v_1, v_2, v_3 be complex numbers linearly independent over $\mathbb{Q}$. Then at least one of the six numbers $e^{u_i v_j}$, $1 \leq i \leq 2, 1 \leq j \leq 3$, is transcendental.*

Status: [classical/open background] — classical, not formalized in this repository in general form.

See Waldschmidt [10, 11]. Every consequence below is tagged [classical/open background] for the same reason: the repository formalizes only the integer-entry specialization of the underlying interpolation determinant, which yields the 3-smooth classification of Section 4 but not the transcendence conclusion.

Lemma 9.4 (Third-base transcendence). *Let $y \in \mathcal{S} \setminus \mathbb{N}_0$ and let $B \in \mathbb{N}_{>0}$ be not three-smooth. Then B^y is transcendental.*

Proof. The numbers $\log 2, \log 3, \log B$ are linearly independent over $\mathbb{Q}$: a nontrivial rational relation among them would give a nontrivial multiplicative relation among 2, 3, and B , which unique factorization forbids because B has a prime factor outside $\{2, 3\}$. Also 1 and y are linearly independent over $\mathbb{Q}$, since y is irrational — a rational solution is integral by the kernel-verified rational-exponent lemma of Section 4, and $y \notin \mathbb{N}_0$.

Apply Theorem 9.3 with $u_1 = 1, u_2 = y$ and $v_1 = \log 2, v_2 = \log 3, v_3 = \log B$. The six values are

$$2, \quad 3, \quad B, \quad 2^y, \quad 3^y, \quad B^y,$$

and the first five are integers, hence algebraic: $2, 3, B$ by hypothesis and $2^y, 3^y$ because $y \in \mathcal{S}$. Therefore the sixth, B^y , is transcendental. $\square$

Status: [classical/open background] — consequence of the six exponentials theorem (Theorem 9.3).

Combining Lemma 9.4 with the kernel-verified 3-smooth classification of Section 4 gives an exact algebraicity trichotomy for integer bases: for every $y \in \mathcal{S} \setminus \mathbb{N}_0$ and every $B \in \mathbb{N}_{>0}$,

$$B^y \in \overline{\mathbb{Q}} \iff B^y \in \mathbb{Z} \iff B = 2^u 3^v \text{ for some } u, v \in \mathbb{N}_0. \quad (11)$$

There is no algebraic-irrational middle case. Any base that is not three-smooth produces a transcendental value outright.

9.3 Products of noninteger solutions

The set $\mathcal{S}$ is closed under addition and under multiplication by natural numbers, but it is not known to be closed under multiplication of two of its own nonintegral elements. The next theorem says exactly what happens at the product exponent xy .

Theorem 9.5 (Product-exponent dichotomy). *Let $x, y \in \mathcal{S}$.*

- (a) *If $x \in \mathbb{N}_0$ or $y \in \mathbb{N}_0$, then $2^{xy} \in \mathbb{N}$ and $3^{xy} \in \mathbb{N}$.*
- (b) *If $x \notin \mathbb{N}_0$ and $y \notin \mathbb{N}_0$, then each of 2^{xy} and 3^{xy} is either an integer or transcendental, at least one of the two is transcendental, and*

$$6^{xy} \text{ is transcendental.}$$

Proof. Part (a) is natural-number dilation: if $y = n \in \mathbb{N}_0$ then $xy = nx$ is a natural multiple of a solution and therefore a solution, and symmetrically in x .

For part (b) put $M_x = 2^x$ and $A_x = 3^x$, both positive integers, and note

$$2^{xy} = M_x^y, \quad 3^{xy} = A_x^y.$$

If M_x is three-smooth, the 3-smooth classification of Section 4 makes M_x^y an integer; if M_x is not three-smooth, Lemma 9.4 makes it transcendental. The same alternative holds for A_x^y . By

Theorem 9.1 applied to the nonintegral solution x , the two outputs M_x and A_x are not both three-smooth, so at least one of $2^{xy}, 3^{xy}$ is transcendental.

Finally $6^x = M_x A_x$ is a positive integer, and it is not three-smooth, since a prime outside $\{2, 3\}$ dividing one factor divides the product. Applying Lemma 9.4 with base 6^x and exponent y gives $(6^x)^y = 6^{xy}$ transcendental. $\square$

Status: [classical/open background] for part (b) — via Theorem 9.3.

Part (a) is elementary and its underlying closure property, that natural multiples of a solution are solutions, is [kernel-verified Lean].

Corollary 9.6 (Diagonal transcendence). *Let $x \in \mathcal{S} \setminus \mathbb{N}_0$. Then 6^{x^2} is transcendental and at least one of $2^{x^2}, 3^{x^2}$ is transcendental. More precisely,*

$$\begin{aligned} 2^{x^2} \in \mathbb{N} &\iff 2^x \text{ is three-smooth,} \\ 3^{x^2} \in \mathbb{N} &\iff 3^x \text{ is three-smooth,} \end{aligned}$$

and failure of either equivalence’s right-hand side makes the corresponding value transcendental.

Proof. Take $y = x$ in Theorem 9.5(b). The two displayed equivalences are the alternative used in that proof, applied to the bases $M_x = 2^x$ and $A_x = 3^x$: a three-smooth base gives an integer by the 3-smooth classification of Section 4, and a base that is not three-smooth gives a transcendental value by Lemma 9.4. $\square$

Corollary 9.7 (Mixed three-smooth bases). *Let $x, y \in \mathcal{S} \setminus \mathbb{N}_0$ and let $u, v \geq 1$ be integers. Then $(2^u 3^v)^{xy}$ is transcendental.*

Proof. The positive integer

$$(2^u 3^v)^x = (2^x)^u (3^x)^v$$

is divisible by every prime occurring in either output, because $u, v \geq 1$. By Theorem 9.1 at least one output has a prime factor outside $\{2, 3\}$, so this base is not three-smooth. Apply Lemma 9.4 with exponent y . $\square$

Status: [classical/open background] for both corollaries — via Theorem 9.3.

9.4 A global second-level trichotomy

The exact generator theorem of Section 5 makes the preceding dichotomy uniform: which of the three patterns occurs is a property of the counterexample as a whole, not of the particular pair (x, y) .

Theorem 9.8 (Uniform second-level pattern). *Assume Conjecture 2.1 fails and let $M = 2^\beta$, $A = 3^\beta$ for the primitive generator β . For every $x, y \in \mathcal{S} \setminus \mathbb{N}_0$, exactly one of the following three patterns holds, and the pattern does not depend on x or y :*

<i>condition on (M, A)</i>	2^{xy}	3^{xy}
<i>M three-smooth, A not</i>	<i>integer</i>	<i>transcendental</i>
<i>A three-smooth, M not</i>	<i>transcendental</i>	<i>integer</i>
<i>neither three-smooth</i>	<i>transcendental</i>	<i>transcendental</i>

Proof. By the exact primitive-generator theorem of Section 5, write $x = n + k\beta$ with $n, k \in \mathbb{N}_0$ and $k \geq 1$, so that

$$2^x = 2^n M^k, \quad 3^x = 3^n A^k.$$

Multiplying by a power of the distinguished smooth prime cannot remove a prime factor outside $\{2, 3\}$, and $k \geq 1$ keeps every such prime of M present in M^k ; hence 2^x is three-smooth exactly when M is, and likewise 3^x is three-smooth exactly when A is. Now apply Theorem 9.5(b) together with the alternative used in its proof. The fourth pattern, both M and A three-smooth, is excluded by Theorem 9.1 applied to β . $\square$

Status: [classical/open background] — via Theorem 9.3.

The structural input, namely $\mathcal{S} = \{n + k\beta : n, k \in \mathbb{N}_0\}$ with unique coordinates, is [kernel-verified Lean] in PrimitiveGenerator (Section 5).

This suggests a clean three-case research split, matching Corollary 9.2. The two one-smooth-output cases reduce to the exceptional primitive cores $w = 3$ and $v = 2$, where the target radicals are 3^ϑ and $2^{1/\vartheta}$ respectively and where Section 6 already locates the wall. In the generic case both second iterates are transcendental and two independent outside-prime supports are available, which is the configuration a mixed archimedean and non-archimedean determinant would exploit; see Section 27.

9.5 Algebraicity formulations equivalent to the conjecture

Theorem 9.9 (Algebraicity tests). *Each of the following is equivalent to Conjecture 2.1.*

- (i) *For every $x \in \mathcal{S}$, $6^{x^2} \in \overline{\mathbb{Q}}$.*
- (ii) *For every $x \in \mathcal{S}$, both $2^{x^2} \in \overline{\mathbb{Q}}$ and $3^{x^2} \in \overline{\mathbb{Q}}$.*
- (iii) *For every $x, y \in \mathcal{S}$, $6^{xy} \in \overline{\mathbb{Q}}$.*
- (iv) *For one fixed pair of integers $u, v \geq 1$, every $x \in \mathcal{S}$ satisfies $(2^u 3^v)^{x^2} \in \overline{\mathbb{Q}}$.*
- (v) *For every $x \in \mathcal{S}$, the integer 6^x is three-smooth.*

Proof. If Conjecture 2.1 holds, then $\mathcal{S} = \mathbb{N}_0$, every exponent appearing in (i)–(iv) is a nonnegative integer, and all the listed values are positive integers, hence algebraic; condition (v) is immediate because $6^x = 2^x 3^x$ is then a power of 6.

Conversely, suppose the conjecture fails and fix $x \in \mathcal{S} \setminus \mathbb{N}_0$. Then (i) and (ii) fail by Corollary 9.6, (iii) fails by taking $y = x$ in Theorem 9.5(b), (iv) fails by Corollary 9.7 with $y = x$, and (v) fails by Theorem 9.1. $\square$

Status: [classical/open background] for (i)–(iv) — via Theorem 9.3. [paper proof in this report] for (v).

Equivalence (v) needs only Theorem 9.1, hence only the kernel-verified transcendence of ϑ ; it is therefore the one item in this list that is within immediate reach of the Lean development.

Items (i)–(iv) replace integrality by the apparently much weaker target of algebraicity, and the six exponentials theorem makes the replacement lossless. That sharpens the statement of what is missing: one would need an independent reason why some nonlinear iterate such as 6^{x^2} is algebraic, derived from the two original integer values 2^x and 3^x alone. No mechanism in the present corpus produces algebraicity at the second level, which is consistent with the four-exponentials diagnosis of Section 7: the two-base configuration supplies only two independent logarithmic directions, and squaring the exponent does not create a third.

9.6 Formalization boundary

The division of labour in this section is unusually clean, and it marks exactly where the repository can advance without new transcendence theory.

Theorem 9.1, Corollary 9.2, part (a) of Theorem 9.5, and equivalence (v) of Theorem 9.9 use only kernel-verified ingredients: transcendence of ϑ , unique factorization, the closure of $\mathcal{S}$ under natural multiples, the simultaneous output normalization, and the 3-smooth classification in `ThreeSmooth`. A proposed module `OutputSmoothness` would host them. With the predicate

$$\text{ThreeSmooth}(m) :\iff \exists u, v \in \mathbb{N}_0, m = 2^u 3^v,$$

the headline statement is a biconditional between integrality of the exponent and joint smoothness of the two outputs:

```
theorem integer_iff_outputs_threeSmooth
  {x : R} (hx : TwoBaseIntegralSolution x) :
  x in Set.range ((^) : Z -> R) <->
    ThreeSmooth (natOutputTwo hx) /\
    ThreeSmooth (natOutputThree hx)
```

Listing 1: Proposed output-smoothness signature.

The reverse direction is the interesting one, and it follows the proof of Theorem 9.1: take logarithms of the two smooth factorizations, assemble (10) as a polynomial of degree at most two in ϑ with integer coefficients, apply

`transcendental_logThreeDivLogTwo`

to force every coefficient to vanish, and discharge the remaining natural-number arithmetic with `omega`. Mathlib’s polynomial evaluation interface is the natural way to build $qX^2 + (p - s)X - r$ and invoke the definition of transcendence.

Everything else in this section — Lemma 9.4, part (b) of Theorem 9.5, Corollaries 9.6 and 9.7, Theorem 9.8, and items (i)–(iv) of Theorem 9.9 — needs the six exponentials theorem in general form. The repository’s interpolation determinant is currently specialized to integer entries: a candidate matrix has integer entries, a nonzero integer determinant satisfies $|D| \geq 1$, and that inequality

is what drives the 3-smooth classification. The natural formal route to Lemma 9.4 is to replace the bound $|D| \geq 1$ by a number-field norm bound for a nonzero *algebraic* determinant, retaining the same analytic upper bound from the interpolation estimate. Packaging that as an “algebraic entries of bounded height” interface would recover the transcendence conclusion while reusing the existing determinant infrastructure, and would place the general six exponentials theorem within the same framework. This is a substantial addition, not a refactor, and it is listed among the transcendence-theoretic priorities of Section 27.

10 Perfect-power content, affine descent, and primitive-output statistics

Section 5 closes the conditional structure question: if Conjecture 2.1 fails, then $\mathcal{S}$ is the free rank-two additive monoid generated by 1 and the canonical least nonintegral solution β , and the least output pair $(M, A) = (2^\beta, 3^\beta)$ is affinely primitive. Both facts are [\[kernel-verified Lean\]](#). What that corpus does not yet record is the *exact* perfect-power content of the output pair at an arbitrary point of the monoid: the kernel-verified corollaries exclude matched decompositions at β , but say nothing about which matched decompositions actually occur elsewhere.

This section supplies that classification. It determines exactly which simultaneous factorizations

$$2^x = 2^r U^d, \quad 3^x = 3^r V^d$$

can occur for $x \in \mathcal{S}$, identifies the maximal common perfect-power degree of the output pair with a gcd of coordinates, and then converts the classification into exact limiting statistics for primitive output pairs. Everything here is conditional on failure and rests on the kernel-verified monoid classification; the classification and the statistics themselves are [\[paper proof in this report\]](#).

Throughout, assume failure and fix the canonical generator β of Section 5, with

$$M = 2^\beta, \quad A = 3^\beta.$$

For $x = n + k\beta$ with $n, k \in \mathbb{N}_0$, exact coordinates give

$$2^x = 2^n M^k, \quad 3^x = 3^n A^k, \tag{12}$$

and the pair (n, k) is uniquely determined by x . The coordinate gcd $\gcd(n, k)$ turns out to be an intrinsic arithmetic invariant of the output pair, not merely of its coordinates.

10.1 Simultaneous perfect powers

Theorem 10.1 (Exact common-power criterion). *Let $x = n + k\beta \in \mathcal{S}$ and let $r \geq 1$. The following are equivalent:*

- (i) *both 2^x and 3^x are perfect r -th powers in $\mathbb{N}$;*
- (ii) *$x/r \in \mathcal{S}$;*
- (iii) *$r \mid n$ and $r \mid k$.*

Proof. If $2^x = U^r$ and $3^x = V^r$ with $U, V > 0$, uniqueness of positive real r -th roots gives

$$2^{x/r} = U, \quad 3^{x/r} = V,$$

so $x/r \in \mathcal{S}$. The converse follows by raising the two integer values at x/r to the r -th power.

For the coordinate form, write $x/r = p + q\beta$ in the unique coordinates supplied by the exact primitive-generator theorem of Section 5. Then $n + k\beta = rp + rq\beta$, and uniqueness of coordinates forces $n = rp$ and $k = rq$, which is exactly the pair of divisibility conditions. Conversely, if $r \mid n$ and $r \mid k$, then $(n/r) + (k/r)\beta \in \mathcal{S}$ and equals x/r . $\square$

For $x > 0$ define the *common output power degree*

$$g(x) := \max\{r \geq 1 : 2^x \text{ and } 3^x \text{ are both perfect } r\text{-th powers}\}.$$

Corollary 10.2 (Coordinate gcd). *If $x = n + k\beta > 0$, then*

$$g(x) = \gcd(n, k),$$

with the convention $\gcd(n, 0) = n$. The output pair is simultaneously power-primitive exactly when $\gcd(n, k) = 1$.

Proof. By Theorem 10.1 the admissible degrees r are precisely the common divisors of n and k , and the set of common divisors of a pair that is not $(0, 0)$ has a largest element, namely $\gcd(n, k)$. $\square$

Consequently every nonzero output pair has a unique decomposition

$$(2^x, 3^x) = (U^g, V^g), \quad g = \gcd(n, k),$$

where $(U, V) = (2^{x/g}, 3^{x/g})$ is a simultaneously power-primitive output pair belonging to the solution x/g . The repository proves power-primitivity for the least pair only; this corollary computes the intrinsic power degree at every point of the solution monoid. For integral $x = n$ it returns $g(n) = n$, which is the familiar statement that 2^n and 3^n are perfect n -th powers and no better; for $x = \beta$ it returns $g(\beta) = \gcd(0, 1) = 1$, which is exactly the kernel-verified exclusion of a common perfect-power degree at the least pair.

Status: [\[paper proof in this report\]](#) — Theorem 10.1 and Corollary 10.2 are paper results of this report. Their inputs are kernel-verified: the unique-coordinate clause of the exact primitive-generator theorem (`PrimitiveGenerator`) and the least-solution corollaries of `LeastSolution`. The implication (i) $\Rightarrow$ (ii) is the special case $r = 0$ of the Lean affine descent theorem `affine_descent_integral_powers` in `ArithmeticRigidity`; the converse and the coordinate criterion (iii) are new here.

10.2 Matched affine decompositions

A descent that strips the same distinguished-base depth and the same outer power from both outputs is the natural common generalization of the two operations above. The repository’s affine descent theorem shows that such a stripping produces a solution; the following criterion says exactly when one exists.

Theorem 10.3 (Matched affine decomposition criterion). *Let $x = n + k\beta \in \mathcal{S}$, let $d \geq 1$, and let $r \in \mathbb{N}_0$. The following are equivalent:*

- (i) *there exist $U, V \in \mathbb{N}_{>0}$ with*

$$2^x = 2^r U^d, \quad 3^x = 3^r V^d;$$
- (ii) *$(x - r)/d \in \mathcal{S}$;*
- (iii) *$0 \leq r \leq n$, $d \mid k$, and $d \mid (n - r)$.*

When these conditions hold,

$$\frac{x - r}{d} = \frac{n - r}{d} + \frac{k}{d}\beta, \quad U = 2^{(x-r)/d}, \quad V = 3^{(x-r)/d}.$$

Proof. Conditions (i) and (ii) are equivalent by uniqueness of positive real roots, applied after removing the common factor 2^r from the first output and 3^r from the second. Expanding (ii) in unique coordinates gives

$$n + k\beta = r + d(p + q\beta)$$

for some $p, q \in \mathbb{N}_0$, so $n = r + dp$ and $k = dq$; this is exactly (iii), and it also yields the displayed formulas. Conversely, (iii) lets one define $p = (n - r)/d \in \mathbb{N}_0$ and $q = k/d \in \mathbb{N}_0$, and then $p + q\beta \in \mathcal{S}$ equals $(x - r)/d$. $\square$

At the least pair the criterion collapses. For $x = \beta$ we have $n = 0$ and $k = 1$, so (iii) forces $r = 0$ and $d = 1$: no nontrivial matched decomposition of (M, A) exists. That is the kernel-verified exact affine primitivity statement of Section 5, recovered here as one instance of a criterion valid at every point.

For a nonintegral solution $x = n + k\beta$ with $k > 0$, Euclidean division $n = qk + r$ with $0 \leq r < k$ gives a canonical reduced affine presentation

$$x = r + k(q + \beta), \tag{13}$$

and hence

$$2^x = 2^r (2^q M)^k, \quad 3^x = 3^r (3^q A)^k.$$

The normalization $0 \leq r < k$ makes the presentation unique. It separates the outer affine degree k — an invariant of the point, being its β -coordinate — from the reduced common base depth r .

Status: [\[paper proof in this report\]](#) — Theorem 10.3 and the reduced presentation (13) are paper results. The implication (i) $\Rightarrow$ (ii) is the kernel-verified affine descent theorem; the equivalence with the coordinate condition (iii), and hence the classification of *all* matched decompositions, is new.

10.3 What is already formal, and what is new here

The table below separates the two layers. The middle column is the kernel-verified statement about the least pair (M, A) ; the right column is what this section adds for a general point $x = n + k\beta$ of the monoid.

Statement	Kernel-verified at β	Extension in this section
Matched base depth	$\min\{v_2(M), v_3(A)\} = 0$, in LeastSolution	Admissible depths at x are exactly $0 \leq r \leq n$ with $d \mid (n - r)$
Common perfect-power degree	no $k \geq 2$ with M, A both k -th powers, in LeastSolution	the exact degree is $g(x) = \gcd(n, k)$ (Corollary 10.2)
Affine primitivity	$M = 2^r U^k$, $A = 3^r V^k$ force $r = 0$, $k = 1$, in LeastSolution	full criterion for every x , with the least pair as the case $(n, k) = (0, 1)$ (Theorem 10.3)
Affine descent	$(x - r)/k \in \mathcal{S}$ from a matched factorization, in ArithmeticRigidity	the descent condition is an equivalence, not merely an implication
Output normalization	$M = 2^a w^d$, $A = 3^b v^e$ with $a = 0$ or $b = 0$ and $\gcd(\gcd(d, e), b - a) = 1$, in OutputNormalization	unchanged; that invariant constrains the two <i>separate</i> degrees, whereas g measures only the common one

The kernel-verified identifiers behind the middle column are

`IsLeastTwoBaseNonintegerSolution.not_both_base_dvd_outputs`

`IsLeastTwoBaseNonintegerSolution.no_common_perfect_power`

`IsLeastTwoBaseNonintegerSolution.affine_primitive`

in **LeastSolution**, together with `exists_simultaneous_primitive_output_normalization` in **OutputNormalization** and `affine_descent_integral_powers` in **ArithmeticRigidity**. None of the right-hand column is machine-checked. A formalization would be light: the coordinate criteria follow from the already accepted unique-representation clause, so the missing work is bookkeeping over $\mathbb{N}_0^2$ rather than new arithmetic.

10.4 Why this does not yet descend below β

If n and k have a common divisor, Theorem 10.1 descends to $x/\gcd(n, k)$. If a matched decomposition exists, Theorem 10.3 descends to $(x - r)/d$. But exact coordinates guarantee that the descended point is again $p + q\beta$ with $p, q \in \mathbb{N}_0$, and it lies below β only when $q = 0$, in which case it is integral and carries no information. The least nonintegral point β is already primitive for every descent of this shape, which is precisely why the kernel-verified corollaries at β are exclusions rather than contradictions.

A successful descent argument therefore needs a transformation not captured by common perfect powers and common distinguished-base depths. Both operations act inside the monoid $\mathbb{N}_0^2$; anything that could close the problem must leave it.

10.5 Total and primitive counts

Corollary 10.2 converts ordinary coprimality statistics of lattice points into exact statistics of hypothetical output pairs. For $T > 0$ put

$$N_\beta(T) = \#\{(n, k) \in \mathbb{N}_0^2 \setminus \{(0, 0)\} : n + k\beta < T\}$$

and

$$P_\beta(T) = \#\{(n, k) \in \mathbb{N}_0^2 : n + k\beta < T, \gcd(n, k) = 1\}.$$

The first counts nonzero solutions below T ; by Corollary 10.2 the second counts exactly those output pairs with no nontrivial common perfect-power degree.

Counting each column separately, for every real $U > 0$,

$$N_\beta(U) = \sum_{0 \leq k < U/\beta} [U - k\beta] - 1 = \frac{U^2}{2\beta} + O_\beta(U + 1), \quad (14)$$

uniformly in U . For integer T this is the kernel-verified cumulative count $C(T)$ of Section 5 minus the point $x = 0$, so (14) agrees in leading order with `QuadraticAsymptotic`; the uniform real-variable form used below is elementary.

Theorem 10.4 (Primitive output pairs). *As $T \rightarrow \infty$,*

$$P_\beta(T) = \frac{T^2}{2\beta\zeta(2)} + O_\beta(T \log T) = \frac{3T^2}{\pi^2\beta} + O_\beta(T \log T).$$

Consequently

$$\frac{P_\beta(T)}{N_\beta(T)} \longrightarrow \frac{1}{\zeta(2)} = \frac{6}{\pi^2}.$$

Proof. Every nonzero coordinate pair has the unique factorization

$$(n, k) = d(n_0, k_0), \quad d = \gcd(n, k), \quad \gcd(n_0, k_0) = 1,$$

and $n + k\beta < T$ holds if and only if $n_0 + k_0\beta < T/d$. Hence

$$N_\beta(T) = \sum_{d \geq 1} P_\beta(T/d),$$

a finite sum for each T , and Möbius inversion gives

$$P_\beta(T) = \sum_{d \leq T} \mu(d) N_\beta(T/d).$$

Inserting the uniform estimate (14),

$$\begin{aligned} P_\beta(T) &= \frac{T^2}{2\beta} \sum_{d \leq T} \frac{\mu(d)}{d^2} + O_\beta \left(\sum_{d \leq T} \left(\frac{T}{d} + 1 \right) \right) \\ &= \frac{T^2}{2\beta\zeta(2)} + O_\beta(T \log T), \end{aligned}$$

where the tail $\sum_{d > T} \mu(d)/d^2 = O(1/T)$ contributes $O_\beta(T)$. Dividing by (14) gives the stated limit. $\square$

The constant $6/\pi^2$ is universal: it does not depend on the hypothetical generator at all. The geometry of the monoid contributes the factor $1/(2\beta)$, and the primitive-lattice density contributes $1/\zeta(2)$. In particular no choice of counterexample can make the primitive proportion small.

10.6 Distribution of the maximal common power degree

For fixed $g \geq 1$ let

$$P_{\beta,g}(T) := \#\{x = n + k\beta : 0 < x < T, \gcd(n, k) = g\}$$

be the number of solutions below T whose output pair has common power degree exactly g .

Corollary 10.5 (Exact-degree distribution). *For fixed $g \geq 1$,*

$$P_{\beta,g}(T) = P_{\beta}(T/g) = \frac{3T^2}{\pi^2 \beta g^2} + O_{\beta}\left(\frac{T}{g} \log \frac{2T}{g}\right).$$

The limiting probability that an output pair below T has maximal common power degree exactly g is

$$\frac{6}{\pi^2 g^2}.$$

Proof. The map $(n_0, k_0) \mapsto g(n_0, k_0)$ is a bijection from the primitive pairs below T/g onto the pairs of gcd exactly g below T , so $P_{\beta,g}(T) = P_{\beta}(T/g)$; now apply Theorem 10.4 and divide by (14). $\square$

The probabilities sum to one, since $\sum_{g \geq 1} 6/(\pi^2 g^2) = 1$. Likewise, by Theorem 10.1 both outputs are perfect r -th powers exactly when $r \mid n$ and $r \mid k$, so

$$\#\{0 < x < T : 2^x \text{ and } 3^x \text{ are both perfect } r\text{-th powers}\} = N_{\beta}(T/r), \quad (15)$$

and the limiting proportion is $1/r^2$.

10.7 Exact generating series

The coordinate monoid has an equally explicit Laplace series. For $s > 0$,

$$\begin{aligned} Z_{\beta}(s) &:= \sum_{x \in \mathcal{S} \setminus \{0\}} e^{-sx} = \sum_{n, k \geq 0} e^{-s(n+k\beta)} - 1 \\ &= \frac{1}{(1 - e^{-s})(1 - e^{-\beta s})} - 1, \end{aligned} \quad (16)$$

the two factors being the two free generators 1 and β . Möbius inversion in the degree gives the primitive-output series

$$Z_{\beta}^{\text{prim}}(s) = \sum_{d \geq 1} \mu(d) Z_{\beta}(ds).$$

At the candidate level the same factorization appears multiplicatively: since $\mathcal{C} = \{2^n M^k : n, k \in \mathbb{N}_0\}$ with unique coordinates,

$$\sum_{m \in \mathcal{C}} m^{-s} = \frac{1}{(1 - 2^{-s})(1 - M^{-s})}, \quad \Re s > 0. \quad (17)$$

These identities make the two free generators visible in analytic form and give clean starting points for weighted Tauberian refinements of Theorem 10.4.

Status: [\[paper proof in this report\]](#) — Theorem 10.4, Corollary 10.5, and the series (16)–(17) are paper results of this report. Their only nonelementary input is the kernel-verified exact monoid classification of Section 5, whose leading-order count is itself [\[kernel-verified Lean\]](#) through `ExactCounting` and `QuadraticAsymptotic`. The finite scans of Section 25 play no role here; the exponent bound in force remains the kernel-verified $x \geq 12$.

10.8 Interpretation

Roughly 60.79% of hypothetical output pairs would be simultaneously power-primitive, and the remaining ones split across degrees $g \geq 2$ with the exact frequencies $6/(\pi^2 g^2)$. This is not pressure toward a contradiction. It is precisely the primitive proportion expected of a free rank-two monoid, and it would be the same for any generator: the statistic is diagnostic, not restrictive.

Its diagnostic value is a warning about method. The kernel-verified primitivity corollaries at β , and the classification of this section at a general point, both say that perfect-power and matched-depth structure is the exception among hypothetical solutions. Any descent argument built on extracting a common root or a common base depth therefore acts on a minority of the conditional population — asymptotically the proportion $1 - 6/\pi^2 \approx 0.3921$ of pairs that admit any nontrivial common degree at all — and leaves the primitive majority untouched. A method that closes the problem must engage the power-primitive pairs directly, which is exactly where the transcendence content of Section 7 sits.

11 Candidate divisibility and ambient root saturation

The exact classification of Section 5 describes the conditional candidate set by its *generators*: under failure, $\mathcal{C} = \{2^n M^k : n, k \in \mathbb{N}_0\}$ with unique coordinates. Generation alone says nothing about how candidates sit inside $\mathbb{N}$ under the two operations that descent arguments and searches actually use — divisibility, and extraction of roots. This section answers both exactly. The candidate monoid has a simple valuation geometry: divisibility is coordinatewise order on a two-dimensional lattice cone, $\mathcal{C}$ is closed under gcd and lcm, and the set of integers some power of which is a candidate is again an explicit cone — one that contains $\mathcal{C}$ with an exact finite-index defect measured by the rational-power index d of Section 5, which is also the exact radical degree of Section 6.

Every statement below is [\[paper proof in this report\]](#). Its inputs — the canonical primitive odd core, the rational-power index, and the exact primitive-generator theorem — are [\[kernel-verified Lean\]](#); what is added here is elementary lattice combinatorics on top of them. None of it is a step toward a contradiction: as the closing discussion makes explicit, the closure properties proved here are compatible with failure by construction, and the kernel-verified exponent bound remains $x \geq 12$.

11.1 A two-dimensional valuation cone

Lemma 11.1 (Valuation coordinates and divisibility). *Assume Conjecture 2.1 is false and let w, d, a, β, M be the canonical data of Section 5, so that $w > 1$ is odd and power-primitive, $d > 0$, and*

$$M = 2^\beta = 2^a w^d.$$

Put $U := w^d$, the odd part of M , and for $n, k \in \mathbb{N}_0$ write

$$m(n, k) := 2^n M^k = 2^{n+ak} U^k, \quad s := v_2(m(n, k)) = n + ak.$$

Then:

- (i) $U > 1$;
- (ii) the map $m(n, k) \mapsto (s, k)$ is a bijection from $\mathcal{C}$ onto the lattice cone

$$\Lambda_a := \{(s, k) \in \mathbb{N}_0^2 : s \geq ak\};$$

- (iii) writing $m_i = m(n_i, k_i)$ and $s_i = n_i + ak_i$,

$$m_1 \mid m_2 \iff s_1 \leq s_2 \text{ and } k_1 \leq k_2. \quad (18)$$

Proof. (i) If $U = 1$ then $M = 2^a$ and $\beta = a \in \mathbb{Z}$, contradicting the irrationality of the least generator.

(ii) By the exact primitive-generator theorem of Section 5, every candidate is $m(n, k)$ for a unique pair $(n, k) \in \mathbb{N}_0^2$. The reparameterization $(n, k) \mapsto (n + ak, k)$ is affine with inverse $(s, k) \mapsto (s - ak, k)$, and $n \geq 0$ holds exactly when $s \geq ak$.

(iii) Since U is odd, $v_2(m(n, k)) = s$, and for an odd prime p we have $v_p(m(n, k)) = k v_p(U)$, which vanishes unless $p \mid U$. Divisibility in $\mathbb{N}$ is comparison of all prime valuations. If $m_1 \mid m_2$ then comparison at 2 gives $s_1 \leq s_2$, and comparison at any prime $p \mid U$ — one exists by (i) — gives $k_1 v_p(U) \leq k_2 v_p(U)$ with $v_p(U) > 0$, hence $k_1 \leq k_2$. Conversely those two inequalities force every prime valuation of m_1 to be at most the corresponding valuation of m_2 . $\square$

Thus candidate divisibility is nothing but the coordinatewise order on Λ_a . The inequality $s \geq ak$ is the only constraint, and it records the base-adic depth a of the least generator: a candidate cannot carry the odd part U^k without carrying at least ak factors of two along with it.

11.2 The candidate gcd/lcm lattice

Theorem 11.2 (Candidate gcd/lcm lattice). *Unconditionally, $\mathcal{C}$ is closed under gcd and lcm. If Conjecture 2.1 fails, then in the notation of Lemma 11.1,*

$$\begin{aligned} \gcd(m_1, m_2) &= m(\min(s_1, s_2) - a \min(k_1, k_2), \min(k_1, k_2)), \\ \text{lcm}(m_1, m_2) &= m(\max(s_1, s_2) - a \max(k_1, k_2), \max(k_1, k_2)); \end{aligned}$$

equivalently, in cone coordinates gcd and lcm are coordinatewise minimum and maximum on Λ_a . Candidate divisibility is therefore a distributive lattice, and Λ_a is a sublattice of $\mathbb{N}_0^2$.

Proof. If the conjecture holds then $\mathcal{C} = \{2^n : n \in \mathbb{N}_0\}$, whose gcds and lcms are again powers of two; this gives the unconditional half.

Assume failure. By Lemma 11.1 the candidate $m(n, k)$ has valuation vector s at the prime 2 and $k v_p(U)$ at each odd p . Gcd and lcm take coordinatewise minima and maxima of valuation vectors, so

$$\gcd(m_1, m_2) = 2^{\min(s_1, s_2)} U^{\min(k_1, k_2)}, \quad \text{lcm}(m_1, m_2) = 2^{\max(s_1, s_2)} U^{\max(k_1, k_2)}.$$

It remains to check that both exponent pairs lie in Λ_a . Suppose $s_1 \leq s_2$; then $\min(s_1, s_2) = s_1 \geq ak_1 \geq a \min(k_1, k_2)$ because $a \geq 0$. For the maximum, suppose $k_1 \leq k_2$; then $\max(s_1, s_2) \geq s_2 \geq ak_2 = a \max(k_1, k_2)$. Hence Λ_a is closed under coordinatewise min and max, that is, it is a sublattice of $\mathbb{N}_0^2$, and the translation back to (n, k) coordinates is the stated substitution $n = s - ak$. A sublattice of a product of chains is distributive. $\square$

Status: [\[kernel-verified Lean\]](#) — `twoBaseNaturalCandidate_gcd` and `twoBaseNaturalCandidate_lcm` in `CandidateLattice`, resting on the coordinatewise arithmetic lemmas `gcd_two_pow_mul_odd_pow` and `lcm_two_pow_mul_odd_pow` of the same module and on the [\[kernel-verified Lean\]](#) exact primitive-generator theorem of Section 5. The formal statements are unconditional: the Lean proof splits on the conjecture and disposes of the affirmative branch through the powers of two. No transcendence input and no analytic estimate enters.

Corollary 11.3 (Counting candidate divisors). *Under failure, the number of candidates dividing $m(n, k)$, including 1 and $m(n, k)$ itself, is*

$$(k+1)(n+1) + \frac{ak(k+1)}{2}.$$

Proof. By (18) a candidate divisor corresponds to a point $(s', j) \in \Lambda_a$ with $0 \leq j \leq k$ and $aj \leq s' \leq n + ak$. For fixed j there are $n + ak - aj + 1$ admissible values of s' . Summing the arithmetic progression over $j = 0, 1, \dots, k$ gives $(k+1)(n+1) + a \sum_{i=0}^k i$, which is the stated value. $\square$

Remark 11.4 (Candidates are not divisor-closed). The full divisor count of $m(n, k) = 2^s U^k$ is $(s+1) \prod_{p|U} (k v_p(U) + 1)$, which for $k \geq 1$ and U with more than one prime factor grows much faster than the count of Corollary 11.3. So $\mathcal{C}$ is emphatically not closed under taking divisors; it is closed only under the two lattice operations. The gap between the two counts is a concrete measure of how thin the candidate sublattice is inside the divisor lattice of a single candidate.

11.3 Fixed prime support

The lattice picture has an immediate consequence for the shape of a hypothetical counterexample branch: the supports do not vary at all.

Corollary 11.5 (Prime-support rigidity). *Assume failure and use the simultaneous primitive output normalization of Section 5,*

$$M = 2^\beta = 2^a w^d, \quad A = 3^\beta = 3^b v^e.$$

Then for $x = n + k\beta$,

$$2^x = 2^{n+ak} w^{dk}, \quad 3^x = 3^{n+bk} v^{ek},$$

so every nonintegral solution — that is, every solution with $k \geq 1$ — has exactly the same odd prime support $\text{supp}(w)$ in its base-2 output, and exactly the same prime support away from 3, namely $\text{supp}(v)$, in its base-3 output. The only variation across the entire nonintegral branch is whether the distinguished primes 2 and 3 occur at all.

Proof. Substituting M and A into $(2^x, 3^x) = (2^n M^k, 3^n A^k)$ gives the displayed factorizations. For $k \geq 1$ the odd part of 2^x is w^{dk} , and $\text{supp}(w^{dk}) = \text{supp}(w)$ because $dk \geq 1$; similarly for v^{ek} . The exponents $n + ak$ and $n + bk$ may vanish, which is exactly the stated freedom at 2 and 3. $\square$

A counterexample therefore locks the whole nonintegral branch onto a single fixed pair of finite prime sets. Computational reconnaissance should exploit this directly, enumerating primitive cores and supports rather than raw integers scanned independently; the finite verifications reported in Section 25 are [\[software computation\]](#) and this observation does not change their status, only the shape of the search that would extend them.

11.4 Ambient root saturation and the index d

Reparameterizing Lemma 11.1 by the exponent of the primitive core w rather than of $U = w^d$ gives the description that governs roots. With $s = v_2(m)$ and t the exponent of w in the odd part of m ,

$$\mathcal{C} = \{2^s w^t : s, t \in \mathbb{N}_0, d \mid t, ds \geq at\}. \quad (19)$$

Indeed $m = 2^n M^k$ has $s = n + ak$ and $t = dk$, so $d \mid t$, and $n \geq 0$ is equivalent to $s \geq at/d$, that is, to $ds \geq at$.

Definition 11.6 (Ambient root saturation). For a multiplicative submonoid $H \subseteq \mathbb{N}_{>0}$ put

$$\text{Sat}(H) := \{u \in \mathbb{N}_{>0} : u^r \in H \text{ for some } r \geq 1\},$$

the set of positive integers admitting some positive power inside H .

Theorem 11.7 (Exact ambient root saturation). *Assume Conjecture 2.1 is false. Then*

$$\text{Sat}(\mathcal{C}) = \{2^s w^t : s, t \in \mathbb{N}_0, ds \geq at\},$$

and for $u = 2^s w^t$ in this set the least positive r with $u^r \in \mathcal{C}$ is

$$r_{\min}(u) = \frac{d}{\gcd(d, t)}.$$

Proof. First let $u = 2^s w^t$ with $s, t \in \mathbb{N}_0$ and $ds \geq at$, and put $r = d/\gcd(d, t)$. Then $rt = d \cdot t/\gcd(d, t)$ is divisible by d , and multiplying $ds \geq at$ by $r > 0$ gives $d(rs) \geq a(rt)$. By (19), $u^r = 2^{rs} w^{rt} \in \mathcal{C}$, so $u \in \text{Sat}(\mathcal{C})$. Conversely, if $u^{r'} \in \mathcal{C}$ for some $r' \geq 1$ then (19) forces $d \mid r't$, since w is power-primitive and hence the exponent of w in the odd part of $u^{r'}$ is exactly $r't$. Now $d \mid r't$ is equivalent to $(d/\gcd(d, t)) \mid r'$, so $r' \geq r$ and r is least.

Now suppose merely that $u \in \mathbb{N}_{>0}$ satisfies $u^r \in \mathcal{C}$ for some $r \geq 1$, and write $u^r = 2^S w^{dk}$ with $dS \geq a dk$. A prime divides u if and only if it divides u^r , so no prime outside $\{2\} \cup \text{supp}(w)$ divides u . Put $c_p := v_p(w)$ for $p \mid w$; comparing valuations in $u^r = 2^S w^{dk}$ gives

$$r v_p(u) = dk c_p \quad \text{for every } p \mid w.$$

Because w is power-primitive, $\gcd\{c_p : p \mid w\} = 1$, so there are integers r_p , almost all zero, with $\sum_p r_p c_p = 1$; then

$$\frac{dk}{r} = \sum_p r_p \cdot \frac{dk c_p}{r} = \sum_p r_p v_p(u) \in \mathbb{Z}.$$

Call this integer $t \geq 0$. Then $v_p(u) = tc_p$ for every $p \mid w$, so the odd part of u is exactly w^t , while $S = rs$ for $s = v_2(u)$. Substituting $S = rs$ and $dk = rt$ into $dS \geq a dk$ and dividing by $r > 0$ yields $ds \geq at$. Hence $u = 2^s w^t$ lies in the stated cone. $\square$

Status: [\[paper proof in this report\]](#) — the only nonformal ingredient is the Bézout step, which is the same primitivity mechanism already kernel-verified in `PrimitiveOddCore` through `oddPrimeValuationGCD`. Everything else is unique factorization and integer arithmetic.

Corollary 11.8 (Ambient saturation index). *Under failure, $\mathcal{C} = \text{Sat}(\mathcal{C})$ if and only if $d = 1$, and in general*

$$d = \min\{r \geq 1 : \exists s \in \mathbb{N}_0, (2^s w)^r \in \mathcal{C}\}.$$

Proof. Take $t = 1$ and any $s \geq a$, so that $ds \geq a = at$; by Theorem 11.7 the integer $2^s w$ lies in $\text{Sat}(\mathcal{C})$ with $r_{\min} = d / \gcd(d, 1) = d$. This proves the displayed formula, and it shows that $2^s w \in \mathcal{C}$ forces $d = 1$. Conversely, if $d = 1$ then the divisibility condition $d \mid t$ in (19) is vacuous and the two sets coincide. $\square$

Corollary 11.9 (Which integer roots of a candidate are again candidates). *Assume failure and let $m = 2^n M^k \in \mathcal{C}$ and $r \geq 1$.*

- (a) *There is an integer u with $u^r = m$ if and only if $r \mid n + ak$ and $r \mid dk$; such a u always lies in $\text{Sat}(\mathcal{C})$.*
- (b) *That integer root is itself a candidate if and only if $r \mid n$ and $r \mid k$, in which case $u = 2^{n/r} M^{k/r}$.*

Proof. (a) Write $m = 2^s w^t$ with $s = n + ak$ and $t = dk$. If $u^r = m$ then, exactly as in the proof of Theorem 11.7, u is supported in $\{2\} \cup \text{supp}(w)$, $r \mid s$, and $t/r \in \mathbb{Z}$ by the Bézout step; conversely $u = 2^{s/r} w^{t/r}$ is an integer r -th root when $r \mid s$ and $r \mid t$. Membership in $\text{Sat}(\mathcal{C})$ is immediate from Definition 11.6.

(b) If $u \in \mathcal{C}$ then $u = 2^{n'} M^{k'}$ for unique $n', k' \in \mathbb{N}_0$, and $u^r = 2^{rn'} M^{rk'}$ together with uniqueness of coordinates gives $n = rn'$ and $k = rk'$. The converse is the direct computation $(2^{n/r} M^{k/r})^r = m$. $\square$

Observation 11.10 (Normal intrinsically, saturated only up to index d). Let $L := \{2^s w^t : s, t \in \mathbb{Z}\}$ be the ambient multiplicative group generated by 2 and w inside $\mathbb{Q}_{>0}^\times$, and let $L_{\mathcal{C}}$ be the group generated by $\mathcal{C}$. Then

$$L_{\mathcal{C}} = \{2^s w^t : s \in \mathbb{Z}, d \mid t\}, \quad [L : L_{\mathcal{C}}] = d, \quad \mathcal{C} = \text{Sat}(\mathcal{C}) \cap L_{\mathcal{C}}.$$

In particular $\text{Sat}(\mathcal{C})$ is itself a submonoid of $\mathbb{N}_{>0}$, namely the full lattice points of the rational cone $ds \geq at, s, t \geq 0$.

Proof. $L_{\mathcal{C}}$ is generated by 2 and $M = 2^a w^d$, so it consists of the elements $2^{i+aj} w^{dj}$ with $i, j \in \mathbb{Z}$, which is the stated set; the quotient $L/L_{\mathcal{C}}$ is isomorphic to $\mathbb{Z}/d\mathbb{Z}$ through the exponent of w . Intersecting the cone description of Theorem 11.7 with the congruence $d \mid t$ returns (19). Closure of the cone under multiplication is additivity of the inequality $ds \geq at$ in (s, t) . $\square$

A terminology caution is worth stating explicitly, because the word “normal” is used differently in affine-semigroup theory. There, normality of a semigroup is measured relative to the group that the semigroup itself generates. By Observation 11.10 that group is $L_{\mathcal{C}}$, in which the congruence $t \equiv 0 \pmod{d}$ already holds identically, and $\mathcal{C}$ is the full set of cone points of $L_{\mathcal{C}}$ for every value of d : intrinsically, $\mathcal{C}$ is normal. Theorem 11.7 measures something different, namely saturation inside the larger ambient lattice L containing every $2^s w^t$. So d is an over-lattice saturation defect of finite index, not a failure of intrinsic normality. This is precisely parallel to its two other exact meanings established elsewhere in this report: d is the least rational-power index of Section 5 and the exact radical degree $[\mathbb{Q}(w^y) : \mathbb{Q}]$ of Section 6. One integer invariant, three characterizations — algebraic degree, least rational power, and lattice saturation defect.

11.5 Why lattice saturation is a good Lean target

Two points must be kept apart. The gcd/lcm closure of Theorem 11.2 and the saturation formula of Theorem 11.7 are compatible with failure by construction: they describe the counterexample world rather than obstruct it, and no contradiction can be extracted from them alone. They are structural information and search normalization. With that said, they are unusually attractive as the next formalization increment.

The mathematical inputs are already kernel-verified. The canonical primitive odd core, the rational-power index, the three-power denominator normalization, and the exact primitive generator with unique coordinates are all [\[kernel-verified Lean\]](#) and live in `ExponentialIdentities/TwoBaseIntegerExponent/`. Everything added in this section is finite combinatorics over those results: unique factorization in $\mathbb{N}$, one Bézout step that already exists in `PrimitiveOddCore` as `oddPrimeValuationGCD`, and elementary min/max arithmetic on $\mathbb{N}_0^2$. There is no analysis, no transcendence input, and no numerical certificate anywhere in the chain — a rare profile in this corpus, and the reason the work should be cheap relative to its yield. A single module, sitting beside `PrimitiveGenerator` and `RationalPowerIndex`, would carry the coordinate bijection of Lemma 11.1, the divisibility criterion (18), the lattice operations, the divisor count, the cone form (19), the saturation set, the exact value of $r_{\min}$, and the index characterization of d . The natural name is

`CandidateLattice`

and the blueprint conventions of Section 28 apply to it unchanged.

The contribution to the search is threefold. First, a kernel-checked normal form: candidate enumeration can be phrased over lattice coordinates rather than over raw integers, so rejection tests become integer inequalities of the form $ds \geq at$ and $d \mid t$ instead of factorizations recomputed for every scanned value. This is exactly the discipline recommended in the computational priority of Section 27, and it applies equally to the certificate architecture that any extension of the [\[software computation\]](#) finite verifications of Section 25 would need; the kernel-verified region itself would still stop at $x \geq 12$ until new certificates are replayed. Second, a single invariant with three independently proved meanings: once Corollary 11.8 is formal alongside the radical-degree proposition of Section 6, any future argument that bounds d — from Kummer theory, from ramification, or from a determinant estimate — bounds all three at once, and the Lean statements make that transfer mechanical rather than a matter of prose. Third, and most concretely, root extraction is the step where descent arguments live. Every descent so far — no matched base depth, no common perfect-power

degree, exact affine primitivity in Section 5 — proceeds by extracting a root of the output pair and contradicting leastness of β . Corollary 11.9 says exactly which integer roots of a candidate survive as candidates and which merely land in the saturation, so a formalized version supplies the reusable side condition that such arguments currently re-prove by hand. That is a modest but genuine improvement in the leverage of the conditional structure, and it costs no new mathematics.

12 A sufficient condition: independence of prime-log products

The reductions of Section 6 isolate the obstruction as a statement about one radical at a time. This section isolates it differently, as a single statement about $\mathbb{Q}$ -linear independence of a natural family of real numbers. The resulting criterion is weaker in appearance than the four exponentials conjecture and is of a different kind: it does not mention exponentials at all.

12.1 Failure as a vanishing linear form in log-products

Suppose the conjecture fails, and let x be a nonintegral solution with outputs

$$M = 2^x \in \mathbb{N}, \quad A = 3^x \in \mathbb{N}.$$

Taking logarithms of $M = 2^x$ and $A = 3^x$ and eliminating x gives the defining relation already recorded in (2):

$$\log M \log 3 - \log A \log 2 = 0. \quad (20)$$

Now expand both integers over their prime factorizations,

$$M = \prod_p p^{\alpha_p}, \quad A = \prod_p p^{\gamma_p}, \quad \alpha_p, \gamma_p \in \mathbb{N}_0,$$

so that $\log M = \sum_p \alpha_p \log p$ and $\log A = \sum_p \gamma_p \log p$. Substituting into (20),

$$\sum_p \alpha_p (\log p \log 3) - \sum_p \gamma_p (\log p \log 2) = 0. \quad (21)$$

Every term of (21) is an integer multiple of a product of exactly two logarithms of primes. Collecting coefficients in the family

$$\{\log p \log q : p \leq q \text{ prime}\},$$

the unordered pair $\{p, 3\}$ with $p \notin \{2, 3\}$ receives the coefficient α_p ; the pair $\{p, 2\}$ with $p \notin \{2, 3\}$ receives $-\gamma_p$; the diagonal pair $\{3, 3\}$ receives α_3 ; the diagonal pair $\{2, 2\}$ receives $-\gamma_2$; and the mixed pair $\{2, 3\}$ receives $\alpha_2 - \gamma_3$.

Thus *failure of the conjecture is exactly the existence of a nontrivial vanishing integer-coefficient linear form in products of two prime logarithms*, with the additional positivity constraint that the coefficients arise from prime factorizations.

12.2 The independence criterion

Definition 12.1 (Prime-log-product independence). Say that *prime-log products are independent* if the real numbers $\log p \log q$, indexed by pairs of primes $p \leq q$, are linearly independent over $\mathbb{Q}$.

Equivalently, in the integer-coefficient form convenient for formalization: for every finite set S of primes and every $c : S \times S \rightarrow \mathbb{Z}$,

$$\sum_{p \in S} \sum_{q \in S} c_{pq} \log p \log q = 0 \quad \implies \quad c_{pq} + c_{qp} = 0 \text{ for all } p, q \in S.$$

(The diagonal case $p = q$ reads $2c_{pp} = 0$, hence $c_{pp} = 0$.) Clearing denominators shows the two formulations agree.

Theorem 12.2 (Prime-log-product independence implies the conjecture). *If prime-log products are independent, then Conjecture 2.1 holds.*

Proof. Let x satisfy $2^x, 3^x \in \mathbb{Z}$. Both powers are positive, so $M := 2^x$ and $A := 3^x$ are positive integers, and $x \geq 0$. Relation (20) holds, hence so does (21).

Apply the independence hypothesis to the finite set $S = \text{primes}(M) \cup \text{primes}(A) \cup \{2, 3\}$ and to the coefficient family

$$c_{pq} = \begin{cases} \alpha_p, & q = 3, \\ -\gamma_p, & q = 2, \\ 0, & \text{otherwise,} \end{cases}$$

which reproduces the left-hand side of (21). The conclusion $c_{pq} + c_{qp} = 0$ yields, coefficient by coefficient:

- taking $q = 3$ and $p \in S \setminus \{2, 3\}$: $\alpha_p + 0 = 0$, so $\alpha_p = 0$;
- taking $p = q = 3$: $2\alpha_3 = 0$, so $\alpha_3 = 0$;
- taking $q = 2$ and $p \in S \setminus \{2, 3\}$: $-\gamma_p + 0 = 0$, so $\gamma_p = 0$;
- taking $p = q = 2$: $-2\gamma_2 = 0$, so $\gamma_2 = 0$;
- taking $p = 2, q = 3$: $\alpha_2 - \gamma_3 = 0$.

Primes outside S divide neither M nor A , so their exponents vanish as well. Hence $\alpha_p = 0$ for every $p \neq 2$, that is

$$M = 2^{\alpha_2},$$

and symmetrically $A = 3^{\gamma_3} = 3^{\alpha_2}$. Writing $n = \alpha_2 \in \mathbb{N}_0$ we get $2^x = M = 2^n$, and strict monotonicity of $t \mapsto 2^t$ gives $x = n \in \mathbb{Z}$. $\square$

Status: [\[kernel-verified Lean\]](#) — `LogProductIndependence` implements this argument, with the independence hypothesis recorded as an explicit `Prop` and never asserted; the accepted conditional theorem is `alaogluErdosConjecture_of_primeLogProductIndependence`.

12.3 What the criterion is, and is not

Three remarks place Theorem 12.2 correctly.

It is a consequence of Schanuel’s conjecture. Schanuel’s conjecture implies that logarithms of multiplicatively independent positive rationals are algebraically independent over $\mathbb{Q}$; in particular

the numbers $\log p$, over all primes p , would be algebraically independent, and then the products $\log p \log q$ are $\mathbb{Q}$ -linearly independent because distinct monomials of degree two in algebraically independent elements are linearly independent. So Definition 12.1 is strictly weaker than Schanuel and lies well inside the standard expectations.

Its strength beyond the conjecture lies in pairs that never occur. This calibration matters and is easy to miss. Collecting the relation (21) back up,

$$\sum_p \alpha_p (\log p \log 3) - \sum_p \gamma_p (\log p \log 2) = \log 3 \log M - \log 2 \log A,$$

so the only products ever exercised are those of the form $\log 2 \log p$ and $\log 3 \log p$. Write F for that sub-family. A vanishing rational relation over F factors as $\log 2 \log u + \log 3 \log v = 0$ for positive rationals u, v , i.e. says precisely that some positive rational $\neq 1$ has a rational ϑ -th power. Equivalently, $\mathbb{Q}$ -linear independence of F is the *rational-output* form of the problem:

$$2^x \in \mathbb{Q} \text{ and } 3^x \in \mathbb{Q} \implies x \in \mathbb{Z}. \quad (22)$$

This is worth separating from Conjecture 2.1. It implies the conjecture, since integers are rational, and it is *a priori strictly stronger*: a rational counterexample such as $2^x = 3/2$ is not an integer counterexample, and nothing in the corpus reduces one to the other. Both follow from four exponentials, because rationals are algebraic.

The proof of Theorem 12.2 establishes (22) verbatim: for rational outputs the prime exponents α_p, γ_p are integers of either sign, and the argument never used their positivity. So the criterion is not a restatement of the conjecture; it delivers the stronger rational-output form. What the calibration does show is where its surplus lies: independence of pairs $\log p \log q$ with $p, q \geq 5$ is never exercised, so one should not expect the full criterion to be easier than (22). Its unconditional dividend is the two-prime case of Section 12.5.

Status: [\[kernel-verified Lean\]](#) for the integer-output form (`alaogluErdosConjecture_of_primeLogProductIndependence`); the rational-output strengthening (22) is [\[paper proof in this report\]](#) and awaits a Lean version stated over `Rat` valuations.

It is nevertheless of a different kind from four exponentials. The four exponentials conjecture is a statement about 2×2 matrices of logarithms and their exponentials; the present criterion mentions no exponentials and is purely a linear-independence assertion about a countable family of real numbers. Neither is known, and neither is known to imply the other in general; what Theorem 12.2 shows is that *either* suffices for the Alaoglu–Erdős conjecture. This is worth recording because the two hypotheses invite different techniques: four exponentials is attacked by auxiliary functions and zero estimates, while linear independence of log-products is a question about the arithmetic of the multiplicative group of $\mathbb{Q}_{>0}$ and its second symmetric power.

The smallest instances are already open. Independence of the family includes, as its very first nontrivial case, the assertion that $(\log 3)^2$, $(\log 2)(\log 3)$ and $(\log 2)^2$ are $\mathbb{Q}$ -linearly independent, equivalently that $\vartheta = \log_2 3$ is not a root of a rational quadratic — a statement about the irrationality of ϑ^2 modulo $\mathbb{Q}(\vartheta)$ that is not decided by present transcendence theory. Indeed the relation exhibited in Section 6 for a 3-smooth counterexample would already contradict the single independence

$$\alpha_2 (\log 2)(\log 3) + \alpha_3 (\log 3)^2 - \gamma_2 (\log 2)^2 - \gamma_3 (\log 2)(\log 3) \neq 0$$

for the relevant nonzero integer vector. Restricted to the two primes 2 and 3, the criterion needed is thus exactly: $1, \vartheta$ and ϑ^2 are linearly independent over $\mathbb{Q}$, which is the 3-smooth shadow of Conjecture 6.3.

12.4 A finite form of the criterion

Because a counterexample has a fixed finite prime support, the full strength of Definition 12.1 is never needed at once.

Corollary 12.3 (Finite prime-support form). *Let S be a finite set of primes containing 2 and 3. If the products $\log p \log q$ with $p \leq q$ in S are $\mathbb{Q}$ -linearly independent, then no counterexample to Conjecture 2.1 has both outputs supported on S .*

Proof. The proof of Theorem 12.2 used only the pairs drawn from the prime support of MA together with 2 and 3. $\square$

12.5 The two-prime case is a theorem, and it is sharp

The finite form is more than a bookkeeping device, because its smallest instance is *provable*.

Proposition 12.4 (Two-prime log-product independence). *Let $p \neq q$ be primes. Then $(\log p)^2$, $\log p \log q$ and $(\log q)^2$ are linearly independent over $\mathbb{Q}$.*

Proof. A vanishing relation $a(\log p)^2 + b \log p \log q + c(\log q)^2 = 0$ with rational a, b, c , divided by $(\log p)^2$, reads $a + bt + ct^2 = 0$ where $t = \log q / \log p$. If $(a, b, c) \neq 0$ this makes t a root of a nonzero rational polynomial, hence algebraic. But t is transcendental: it is irrational by unique factorization, and were it algebraic irrational then Gelfond–Schneider applied to $p^t = q$ would make the integer q transcendental. $\square$

Consequently the hypothesis of Corollary 12.3 is *unconditionally true* for $S = \{2, 3\}$, and the corollary becomes an unconditional exclusion. Stated directly, and proved in Lean without any appeal to Definition 12.1:

Theorem 12.5 (Three-smooth outputs force an integral exponent). *If $2^x \in \mathbb{Z}$ and $3^x \in \mathbb{Z}$ are both 3-smooth — divisible by no prime outside $\{2, 3\}$ — then $x \in \mathbb{Z}$.*

Proof. Write $2^x = 2^a 3^b$ and $3^x = 2^c 3^d$. Taking base-two logarithms gives the two relations $x = a + b\vartheta$ and $x\vartheta = c + d\vartheta$. Substituting the first into the second eliminates x and leaves

$$b\vartheta^2 + (a - d)\vartheta - c = 0.$$

If $b \neq 0$ this makes ϑ a root of a nonzero rational quadratic, contradicting its transcendence. Hence $b = 0$ and $x = a$. $\square$

Corollary 12.6 (A smooth candidate has a rough partner). *If x is a nonintegral solution and 2^x is 3-smooth, then 3^x has a prime factor at least 5. The two outputs of a counterexample can never both be built from $\{2, 3\}$.*

Status: [\[kernel-verified Lean\]](#) — `integer_of_threeSmooth_outputs` and `exists_prime_five_le_dvd_of_threeSmooth_candidate` in `SmoothOutputs`. The proofs use only transcendence of ϑ , itself kernel-verified in this corpus.

Two remarks fix the position of this result precisely.

The two-sided hypothesis is essential. Assuming only that 2^x is 3-smooth gives $x = a + b\vartheta$ and then requires $(3^\vartheta)^b \in \mathbb{Q}$, which is exactly the open localized-radical condition of Section 6. What the second smoothness hypothesis supplies is a *second* linear relation between x and ϑ ; two such relations eliminate x and leave an algebraic equation in ϑ alone. This degree drop is the whole mechanism, and it is worth isolating as a template: *each independent multiplicative constraint on an output buys one linear relation, and two relations force algebraicity of ϑ .* The difficulty of the general problem is precisely that one never gets a second constraint for free.

Three primes is where the criterion becomes conjectural. Proposition 12.4 exhausts what transcendence of a single logarithmic ratio can give. For $S = \{2, 3, 5\}$ a vanishing relation among the six products, divided by $(\log 2)^2$, reads

$$a + b\vartheta + c\vartheta_5 + d\vartheta^2 + e\vartheta\vartheta_5 + f\vartheta_5^2 = 0, \quad \vartheta_5 = \log_2 5,$$

that is, the point (ϑ, ϑ_5) lies on a rational conic. Both coordinates are transcendental and $1, \vartheta, \vartheta_5$ are $\mathbb{Q}$ -independent by unique factorization, but no present method excludes a quadratic relation between them. So Definition 12.1 is a theorem for two primes, open from three onwards, and the Alaoglu–Erdős conjecture sits exactly on that boundary.

13 The algebraic auxiliary-base spectrum

Section 9 classified the *natural* auxiliary bases whose x -th powers remain algebraic, and had to tag the transcendence half of that classification as classical background: the repository’s interpolation determinant carried integer entries, so it could prove $B^x \notin \mathbb{Q}$ but not $B^x \notin \overline{\mathbb{Q}}$. That boundary has moved. The same square determinant now accepts algebraic entries of bounded height inside a number field, and the resulting classification is exact over *all* positive real algebraic bases. A hypothetical nonintegral solution turns out to have one fixed, explicitly described set of good bases — a dense but very thin subgroup of $\mathbb{R}_{>0}$ — and every positive algebraic base outside it produces a transcendental value.

13.1 The saturated group of positive two-three units

The multiplicative group $\overline{\mathbb{Q}}_{>0}$ of positive real algebraic numbers is a vector space over $\mathbb{Q}$ once rational scalars act through the unique positive real root:

$$q \cdot \alpha = \alpha^q, \quad q \in \mathbb{Q}, \quad \alpha \in \overline{\mathbb{Q}}_{>0}.$$

For multiplicatively independent $a, b \in \overline{\mathbb{Q}}_{>0}$ define the *saturated group*

$$\Gamma(a, b) = \{a^r b^s : r, s \in \mathbb{Q}\}.$$

The map $\mathbb{Q}^2 \rightarrow \Gamma(a, b)$, $(r, s) \mapsto a^r b^s$, is an isomorphism of $\mathbb{Q}$ -vector spaces; injectivity is exactly multiplicative independence after clearing denominators. For the original problem write

$$\Gamma = \Gamma(2, 3) = \{2^r 3^s : r, s \in \mathbb{Q}\}.$$

This group has a second description that is the one the Lean development actually uses. Let $\langle 2, 3 \rangle_{\mathbb{Z}} = \{2^i 3^j / (2^k 3^\ell) : i, j, k, \ell \in \mathbb{N}_0\}$ be the group of positive rational 2, 3-units, and let

$$\sqrt[n]{\langle 2, 3 \rangle_{\mathbb{Z}}} := \{\gamma > 0 : \exists n \geq 1, \gamma^n \in \langle 2, 3 \rangle_{\mathbb{Z}}\}$$

be its divisible hull inside $\mathbb{R}_{>0}$.

Observation 13.1. $\sqrt[n]{\langle 2, 3 \rangle_{\mathbb{Z}}} = \Gamma$. In particular every element of the divisible hull is automatically algebraic.

Indeed, if $\gamma^n = 2^i 3^j / (2^k 3^\ell)$ then γ is the unique positive real n -th root of a positive rational, hence $\gamma = 2^{(i-k)/n} 3^{(j-\ell)/n} \in \Gamma$; conversely $2^r 3^s$ raised to a common denominator of r, s lands in $\langle 2, 3 \rangle_{\mathbb{Z}}$. The Lean predicate `HasPositiveTwoThreeUnitPower` is the literal transcription of membership in the hull, so it and membership in Γ may be used interchangeably below.

13.2 The exact spectrum theorem

The kernel-verified corpus contains the determinant in fully symmetric form: three positive algebraic bases with injective nonnegative monomials cannot all have algebraic x -th powers when x is irrational. The next theorem is the coordinate-free synthesis for an arbitrary independent pair; its proof is a direct application of the same classical six-exponentials input, and it is stated here at paper level because the Lean development fixes the pair $(2, 3)$.

Theorem 13.2 (Exact algebraic-base spectrum). *Let $a, b \in \overline{\mathbb{Q}}_{>0}$ be multiplicatively independent, let $x \in \mathbb{R} \setminus \mathbb{Q}$, and assume $a^x, b^x \in \overline{\mathbb{Q}}$. Then, for every $\alpha \in \overline{\mathbb{Q}}_{>0}$,*

$$\alpha^x \in \overline{\mathbb{Q}} \iff \alpha \in \Gamma(a, b).$$

[paper proof in this report]

Proof. If $\alpha = a^r b^s$ with $r, s \in \mathbb{Q}$, positivity gives $\alpha^x = (a^x)^r (b^x)^s$, which is algebraic.

Conversely, suppose $\alpha \notin \Gamma(a, b)$. Then $\log a, \log b, \log \alpha$ are linearly independent over $\mathbb{Q}$: a rational relation clears to an integer multiplicative relation, and if the coefficient of $\log \alpha$ vanishes it contradicts the multiplicative independence of a, b , while otherwise it places α in the saturated group. The pair $1, x$ is linearly independent over $\mathbb{Q}$ because x is irrational. Apply the six exponentials theorem (Theorem 9.3) to the three rows $\log a, \log b, \log \alpha$ and the two columns $1, x$. Five of the six exponential values,

$$a, \quad b, \quad \alpha, \quad a^x, \quad b^x,$$

are algebraic. Therefore the sixth, α^x , is transcendental. $\square$

Conceptually this says more than any list of good auxiliary bases: two algebraic outputs have already exhausted the maximum algebraic multiplicative rank that six exponentials permits.

Specializing to $a = 2, b = 3$ gives the statement the repository proves outright.

Corollary 13.3 (Complete algebraic spectrum of a counterexample). *Let $x \in \mathcal{S} \setminus \mathbb{Z}$. For every $\alpha \in \overline{\mathbb{Q}}_{>0}$,*

$$\alpha^x \in \overline{\mathbb{Q}} \iff \alpha \in \Gamma.$$

In particular the spectrum does not depend on which hypothetical nonintegral solution is chosen.
[\[kernel-verified Lean\]](#)

A rational solution is integral, so x is irrational and Theorem 13.2 applies; the formal route is the divisible-hull equivalence of the next subsection together with Observation 13.1.

Status: [\[kernel-verified Lean\]](#) — the divisible-hull equivalence lives in `AlgebraicBaseUnitPower`.

Its two forms are

`real_rpow_isAlgebraic_iff_integer_or_hasPositiveTwoThreeUnitPower`

together with its counterexample-relative form

`real_rpow_isAlgebraic_iff_hasPositiveTwoThreeUnitPower`

in the `TwoBaseNonintegerSolution` namespace, both resting on the injectivity criterion `real_two_three_gamma_monomial_injective_iff`. The natural-base and rational-base specializations are `rpow_isAlgebraic_iff_integer_or_eq_two_pow_mul_three_pow` in `AlgebraicThirdBase` and `rat_rpow_isAlgebraic_iff_isTwoThreeUnit_of_not_integer` in `AlgebraicRationalBase`; the half-plane, mixed-iterate and rank-one ratio consequences are in `AlgebraicOutputLocus`; and the simultaneous exclusion of both canonical radicals from the entire hull is in `CanonicalRadicalDivisibleHull`.

The group Γ is countable, divisible, of rational rank two, and dense in $\mathbb{R}_{>0}$ — density already follows from the subgroup $\{2^r : r \in \mathbb{Q}\}$. A counterexample would therefore possess a dense set of algebraic bases with algebraic output, while every algebraic base outside that exact dense subgroup would give a transcendental one.

Corollary 13.4 (Rank bound for any family of algebraic outputs). *Let $x \in \mathcal{S} \setminus \mathbb{Z}$ and let $\alpha_1, \dots, \alpha_t \in \overline{\mathbb{Q}}_{>0}$ satisfy $\alpha_i^x \in \overline{\mathbb{Q}}$. Then every α_i lies in Γ and the multiplicative rank of $\alpha_1, \dots, \alpha_t$ is at most two.* [\[paper proof in this report\]](#)

Both ingredients are kernel-verified — membership in the hull by `AlgebraicBaseUnitPower` and the rank ceiling by `AlgebraicRankThreeEquivalence` — but the combined family statement is recorded here at paper level.

13.3 The divisible-hull formulation

Stated without reference to a fixed counterexample, the kernel-verified form is a single biconditional valid for every positive real algebraic base. For $\gamma \in \overline{\mathbb{Q}}_{>0}$, under $2^x, 3^x \in \mathbb{Z}$,

$$\gamma^x \in \overline{\mathbb{Q}} \iff x \in \mathbb{Z} \quad \text{or} \quad \gamma \in \sqrt[\infty]{\langle 2, 3 \rangle_{\mathbb{Z}}}. \tag{23}$$

The arithmetic core is an exact combinatorial equivalence: γ has *no* positive natural power in $\langle 2, 3 \rangle_{\mathbb{Z}}$ precisely when the monomial map $(i, j, k) \mapsto 2^i 3^j \gamma^k$ is injective on $\mathbb{N}_0^3$. The forward implication of

(23) then feeds that injectivity into the generic algebraic determinant, and the reverse implication descends algebraicity from a positive integral power. Specializing γ to a natural number recovers the 3-smooth classification of Section 9 with “rational” upgraded to “algebraic”, so for a hypothetical nonintegral solution the three properties $a^x \in \mathbb{Z}$, $a^x \in \mathbb{Q}$ and $a^x \in \overline{\mathbb{Q}}$ coincide for $a \in \mathbb{N}_{>0}$, and hold exactly at 2, 3-smooth a . Specializing to a positive rational q gives $q^x \in \overline{\mathbb{Q}} \iff q \in \langle 2, 3 \rangle_{\mathbb{Z}}$; there is no algebraic-irrational middle case anywhere in the picture.

13.4 The determinant engines behind the spectrum

Two generic determinant theorems, neither of which mentions the bases 2, 3, carry the whole classification. The rational one states that if $a, b, c \in \mathbb{Q}_{>0}$ have injective nonnegative monomials $(i, j, k) \mapsto a^i b^j c^k$ and $a^x, b^x, c^x \in \mathbb{Q}$, then $x \in \mathbb{Q}$. The algebraic one removes both restrictions: if a, b, c are positive real algebraic numbers with injective nonnegative monomials and a^x, b^x, c^x are all algebraic, then $x \in \mathbb{Q}$. Both allow bases below 1 and negative exponents, and both use $1 + |x|$ in the size estimate.

Status: [\[kernel-verified Lean\]](#) —

`rational_of_three_rat_rpows_rational_of_monomial_injective` in
`RationalSixExponentials`, and
`rational_of_three_real_rpows_isAlgebraic_of_monomial_injective` in
`AlgebraicSixExponentials`.

The step from rational to algebraic outputs is exactly the replacement of the trivial bound $|D| \geq 1$ for a nonzero integer determinant by a field-norm bound for a nonzero algebraic integer determinant, which is the extension anticipated in Section 9. The six input and output values are placed in one number field K with one common nonzero natural scale s making every entry lie in $\mathcal{O}_K$. For side parameter n the square matrix has size n^6 and each entry acquires the factor s^{n^5} , so the determinant acquires $s^{n^{11}}$; at every embedding other than the distinguished one the six factor groups in an entry are bounded by T^{6n^5} , and the norm of the nonzero algebraic-integer determinant gives

$$\frac{1}{((n^6)! T^{6n^{11}})^{[K:\mathbb{Q}]-1}} \leq |\sigma(\det A)|.$$

The verified estimates $(n^6)! \leq 2^{n^{11}}$ and $((n^6)! T^{6n^{11}})^{[K:\mathbb{Q}]-1} \leq ((2T^6)^{[K:\mathbb{Q}]-1})^{n^{11}}$ put the entire conjugate and degree loss on the same n^{11} scale as the old rational denominator loss, where the existing explicit analytic margin already absorbs it. No new analytic estimate was required — only honest bookkeeping.

Status: [\[kernel-verified Lean\]](#) — `one_div_conjugate_entry_bound_le_norm_det` and `one_div_entry_bound_le_scaled_norm_det` in `AlgebraicIntegerDeterminant`, with the number-field realizations `AlgebraicOutputBridge.ofRealIsAlgebraic` and `FiniteAlgebraicOutputBridge.ofFinSixIsAlgebraic`.

13.5 The strong spectrum via Baker and Roy

Let $\mathcal{L}$ denote the $\overline{\mathbb{Q}}$ -vector space generated by 1 and by all logarithms of nonzero algebraic numbers, with arbitrary complex branches allowed; for positive real algebraic numbers we use the real logarithm, which is one of those branches. Roy’s strong six exponentials theorem [8, 11] states that if three complex numbers are linearly independent over $\overline{\mathbb{Q}}$ and two more are linearly independent over $\overline{\mathbb{Q}}$, then at least one of the six pairwise products lies outside $\mathcal{L}$.

Theorem 13.5 (Strong algebraic-base spectrum). *Under the hypotheses of Theorem 13.2, if $\alpha \notin \Gamma(a, b)$, then*

$$x \log \alpha \notin \mathcal{L}.$$

In particular α^x is transcendental. [paper proof in this report]

Proof. As in Theorem 13.2, the logarithms $\log a, \log b, \log \alpha$ are $\mathbb{Q}$ -linearly independent, and Baker’s homogeneous theorem [3] upgrades that to linear independence over $\overline{\mathbb{Q}}$. The number x is transcendental: were it algebraic, its irrationality together with $a \neq 0, 1$ would make a^x transcendental by Gelfond–Schneider. Hence $1, x$ are linearly independent over $\overline{\mathbb{Q}}$. Apply Roy’s theorem to the rows $\log a, \log b, \log \alpha$ and the columns $1, x$. Five products lie in $\mathcal{L}$, namely $\log a, \log b, \log \alpha, x \log a = \log(a^x)$ and $x \log b = \log(b^x)$. Therefore the sixth product $x \log \alpha$ lies outside $\mathcal{L}$. $\square$

Definition 13.6. Call a positive exponential value e^z *log-strongly transcendental* when $z \notin \mathcal{L}$. This implies ordinary transcendence of e^z , since the logarithm of a positive algebraic number belongs to $\mathcal{L}$.

So outside Γ the failure is not merely that α^x is transcendental: the number $x \log \alpha$ is not even a $\overline{\mathbb{Q}}$ -linear combination of 1 and logarithms of algebraic numbers. This is a genuinely stronger exclusion than Corollary 13.3, and it is the form in which the spectrum interacts with linear forms in logarithms.

13.6 One base is enough: detectors and rank equivalences

Because the spectrum is the same for every hypothetical nonintegral solution, a *single* auxiliary base outside Γ already decides the whole conjecture.

Theorem 13.7 (Algebraic detector equivalence). *Fix $\alpha \in \overline{\mathbb{Q}}_{>0} \setminus \Gamma$. Then the Alaoglu–Erdős conjecture is equivalent to*

$$\forall x \in \mathbb{R}, \quad 2^x, 3^x \in \mathbb{Z} \implies \alpha^x \in \overline{\mathbb{Q}},$$

and equally to the assertion that every such x satisfies $x \log \alpha \in \mathcal{L}$. [paper proof in this report]

Proof. If the conjecture holds, every solution is a nonnegative integer and α^x is algebraic, with $x \log \alpha \in \mathcal{L}$. Conversely a nonintegral solution would make α^x transcendental by Corollary 13.3 and would put $x \log \alpha$ outside $\mathcal{L}$ by Theorem 13.5. $\square$

This is a paper-level universalization of a pointwise locus that the kernel already knows. Valid detectors include, for instance,

$$5, \quad \sqrt{5}, \quad 1 + \sqrt{2}, \quad \frac{7 + 3\sqrt{5}}{2},$$

provided the chosen number does not accidentally lie in Γ ; for an algebraic integer with a prime divisor outside $\{2, 3\}$ that exclusion is immediate, and in general it is a multiplicative independence check. Nothing arithmetically special about 5 is used.

Four instances of this phenomenon are themselves kernel-verified, and they are the most quotable consequences of the whole algebraic-output layer.

One-base equivalents of the conjecture. Each of the following is equivalent to the Alaoglu–Erdős conjecture, and each equivalence is accepted by the Lean kernel.

- (i) Every x with $2^x, 3^x \in \mathbb{Z}$ has 5^x algebraic.
- (ii) Every x with $2^x, 3^x \in \mathbb{Z}$ has 6^{x^2} algebraic.
- (iii) For every x with $2^x, 3^x \in \mathbb{Z}$, the positive-algebraic output locus

$$\{\alpha \in \overline{\mathbb{Q}}_{>0} : \alpha^x \in \overline{\mathbb{Q}}\}$$

has multiplicative rank three, i.e. contains a, b, c with injective nonnegative monomials.

- (iv) For every x with $2^x, 3^x \in \mathbb{Z}$, the rational squared-output ratio locus has rank two, i.e. contains algebraic members $(2^i/3^j)^{x^2}$ and $(2^k/3^\ell)^{x^2}$ with $il \neq kj$.

A counterexample would collapse (iii) to rank at most two and confine the entire locus in (iv) to a single rational ray in the exponent lattice. [\[kernel-verified Lean\]](#)

Status: [\[kernel-verified Lean\]](#) — all four equivalences are accepted by the kernel; the Lean names follow.

Items (i) and (ii) are `alaogluErdosConjecture_iff_five_rpow_isAlgebraic` and `alaogluErdosConjecture_iff_six_squared_exponent_isAlgebraic` in `AlgebraicConsequences`, which also carries the two-iterate variant `alaogluErdosConjecture_iff_both_iterated_outputs_isAlgebraic`. Item (iii) is

`alaogluErdosConjecture_iff_algebraicRpowOutputLocusHasRankThree`

in `AlgebraicRankThreeEquivalence`, and item (iv) is

`alaogluErdosConjecture_iff_ratioSquaredAlgebraicLocusHasRankTwo`

in `AlgebraicOutputLocus`. The same modules also record the pointwise forms `ratioSquaredAlgebraicLocusHasRankTwo_iff_integer` and `algebraicRpowOutputLocusHasRankThree_iff_integer`, and the explicit p -adic line `exists_external_prime_algebraic_output_ratio_padic_line` on which a counterexample would force every algebraic ratio output to sit.

13.7 Exact exponent, rational-function, and tower loci

Fix a hypothetical nonintegral solution x and put

$$\mathcal{E}_x = \{y \in \mathbb{R} : 2^y, 3^y \in \overline{\mathbb{Q}}\}.$$

The transposed six-exponentials determinant determines this set exactly:

$$\mathcal{E}_x = \mathbb{Q} + \mathbb{Q}x. \quad (24)$$

Thus every exponent outside the affine plane makes at least one of 2^y and 3^y transcendental. In particular $y = x^n$ for every $n \geq 2$, and also $y = x^{-1}$, lie outside the plane because x is transcendental. For every fixed $n \geq 2$, requiring both 2^{x^n} and 3^{x^n} to be algebraic at every original two-base solution is therefore equivalent to Conjecture 2.1. The theorem says “at least one”: it neither selects which base is transcendental nor supplies the missing algebraicity needed to prove the conjecture.

The exact plane makes every genuinely nonaffine rational function an equally sharp detector. For $P, Q \in \mathbb{Q}[X]$ with $Q(x) \neq 0$,

$$2^{P(x)/Q(x)}, 3^{P(x)/Q(x)} \in \overline{\mathbb{Q}} \iff P = (q + rX)Q \text{ for some } q, r \in \mathbb{Q}. \quad (25)$$

The reverse direction is literal affine membership in (24); the forward direction uses transcendence of x to turn equality after evaluation into a polynomial identity. Consequently any fixed nonaffine rational function represented as P/Q with the displayed denominator Q having no real zero gives a global formulation equivalent to Conjecture 2.1, and every $P \in \mathbb{Q}[X]$ of degree at least two does so with $Q = 1$. These are detector equivalences, not new algebraicity results: they relocate the open conclusion to another exponent without proving it.

Under failure, the same module chooses the canonical generator β so that the two loci can be compared on one rational-function chart. For every P/Q defined at β , both $2^{P(\beta)/Q(\beta)}$ and $3^{P(\beta)/Q(\beta)}$ are algebraic exactly when $P = (q + rX)Q$ with $q, r \in \mathbb{Q}$, whereas $P(\beta)/Q(\beta)$ is itself an original integral solution exactly when $P = (n + kX)Q$ with $n, k \in \mathbb{N}_0$. Thus the exact solution monoid is the nonnegative integral cone inside the rational affine algebraic-output plane; the comparison is exact, but it does not collapse the larger plane to that cone.

There is also an exact two-dimensional picture for the squared monomials. Under a counterexample, the set

$$\mathcal{F}_x = \{(i, j) \in \mathbb{N}_0^2 : (2^i 3^j)^{x^2} \in \overline{\mathbb{Q}}\}$$

is exactly one of

$$\mathbb{N}_0 \times \{0\}, \quad \{0\} \times \mathbb{N}_0, \quad \{(0, 0)\}.$$

Writing $M = 2^x$ and $B = 3^x$, the three cases say respectively that M is 2, 3-smooth and B is not, that B is 2, 3-smooth and M is not, or that neither is; both cannot be 2, 3-smooth for a nonintegral solution. Under the original two integrality hypotheses this coordinate-face property is pointwise equivalent to nonintegrality, and the conjecture is equivalent to excluding it for every solution. All three faces remain compatible with the present theory, so the trichotomy does not itself yield a contradiction.

A related transition theorem has a different, solution-independent plane. With $\vartheta = \log_2 3$, define

$$\mathcal{T}_x = \{y \in \mathbb{R} : 2^y \in \overline{\mathbb{Q}} \text{ and } 2^{xy} \in \overline{\mathbb{Q}}\}.$$

For every hypothetical nonintegral solution,

$$\mathcal{T}_x = \mathbb{Q} + \mathbb{Q}\vartheta. \quad (26)$$

Two consecutive algebraic base-two tower levels therefore correspond exactly to one exponent in this fixed plane. More strongly, at most one positive depth n can satisfy $x^n \in \mathcal{T}_x$. Indeed, transcendence

of x makes the ϑ -coefficients in the two affine expressions nonzero; raising the two expressions to cross-powers and using transcendence of ϑ promotes equality at ϑ to a polynomial identity, whose degree comparison forces the two depths to coincide. Hence every three consecutive positive-depth levels

$$2^{x^n}, \quad 2^{x^{n+1}}, \quad 2^{x^{n+2}}$$

contain a transcendental value, and over any finite set the collection of depths starting an adjacent algebraic pair has cardinality at most one. Conversely, for any fixed distinct positive depths $n \neq m$, the assertion that every original two-base solution has algebraic adjacent pairs at both depths n and m is equivalent to Conjecture 2.1. The restriction still permits isolated algebraic levels and a single adjacent algebraic pair; it does not prove the full conjecture.

The symmetric transition picture is stronger than merely repeating the same argument with base three. Put

$$\mathcal{S}_x = \{y \in \mathbb{R} : 3^y \in \overline{\mathbb{Q}} \text{ and } 3^{xy} \in \overline{\mathbb{Q}}\}.$$

Under a nonintegral solution,

$$y \in \mathcal{S}_x \iff \vartheta y \in \mathbb{Q} + \mathbb{Q}\vartheta. \tag{27}$$

No positive depth in the base-two tower can start an algebraic adjacent pair simultaneously with any positive depth in the base-three tower, even when the two depths are unrelated. The base-three tower itself has at most one such positive depth. Consequently, across the two principal towers together there is at most one positive transition depth, and at that depth at most one of the two bases can occur. There is also a direction-uniform version in the multiplicatively dependent-output branch: once its kernel relation produces a nonzero rational 2, 3-unit direction with a depth-one transition, every nonzero rational 2, 3-unit direction is excluded at every depth $n \geq 2$. These are strict structural obstructions, but they still allow the single depth-one transition forced by the kernel relation and therefore do not close that branch.

Status: [\[kernel-verified Lean\]](#) — all statements in this subsection are accepted by the kernel.

The exact exponent plane and its natural-power and reciprocal consequences are `TwoBaseNonintegerSolution.twoThreeRpowAlgebraic_iff_affine`, `transcendental_two_or_three_rpow_nat_pow`, `transcendental_two_or_three_rpow_inv`, and `alaogluErdosConjecture_iff_twoThreeRpowAlgebraic_nat_pow` in `AlgebraicExponentLocus`. Equation (25) and its global and polynomial specializations are `TwoBaseNonintegerSolution.twoThreeRpowAlgebraic_rationalFunction_iff`, `alaogluErdosConjecture_iff_twoThreeRpowAlgebraic_rationalFunction`, and `alaogluErdosConjecture_iff_twoThreeRpowAlgebraic_polynomial_of_two_le_natDegree` in `RationalFunctionAlgebraicLocus`. Its simultaneous two-locus package is `exists_generator_with_exact_rationalFunction_algebraic_and_integral_loci`. The exact squared-coordinate-face statements are `TwoBaseNonintegerSolution.exists_exact_algebraic_output_monomial_locus`, `algebraic_two_three_squared_monomial_locus_trichotomy`, `squaredTwoThreeMonomialAlgebraicLocusIsCoordinateFace_iff_not_integer`, and `alaogluErdosConjecture_iff_no_squaredTwoThreeMonomial_coordinateFace` in `AlgebraicOutputLocus`. Finally, (26), global uniqueness of a transition depth, and the fixed-distinct-depth equivalence are `TwoBaseNonintegerSolution.twoRpowAlgebraicTransition_iff_affine_logRatio`, `transcendental_one_of_three_consecutive_two_rpow_tower`, `not_two_distinct_algebraic_tower_transitions`, `card_algebraicTowerTransitionDepths_le_one`, and `alaogluErdosConjecture_iff_two_distinct_tower_transitions_algebraic` in `AlgebraicTowerTransition`. The symmetric locus (27), the cross-base exclusion, the combined uniqueness theorem, and the rational-unit-direction exclusions are respectively `TwoBaseNonintegerSolution.threeRpowAlgebraicTransition_iff_scaled_affine_logRatio`, `not_baseTwo_and_baseThree_tower_transitions`, `eq_of_twoOrThree_tower_transitions`, and `not_higher_unitDirectionTowerTransition_of_kernelRelation` in `CrossBaseTowerTransition`. Their axiom audit uses only `propext`, `Classical.choice`, and `Quot.sound`. None asserts Conjecture 2.1.

There is a corresponding zero-one law: assuming the conjecture fails and fixing $\alpha \in \overline{\mathbb{Q}}_{>0}$, exactly one of two things happens. Either $\alpha \in \Gamma$ and α^x is algebraic for *every* nonintegral solution x , or $\alpha \notin \Gamma$ and α^x is transcendental for every nonintegral solution. No base can behave one way for one counterexample and the other way for another.

13.8 Why the exact spectrum does not settle the conjecture

It is worth being precise about what has and has not been gained. The spectrum is an exact answer to the question “which algebraic bases stay algebraic?”, and its answer is $\Gamma = \{2^r 3^s : r, s \in \mathbb{Q}\}$ — but that is precisely the set of bases whose logarithms lie in the rational span of $\log 2$ and $\log 3$. Every base the spectrum certifies is, logarithmically, a rational recombination of the two columns already present in the four-exponentials matrix of Section 7. Six exponentials needs a

third row genuinely independent of $\log 2, \log 3$ together with algebraicity of its output; the spectrum theorem shows that a counterexample supplies algebraic outputs at exactly the bases that fail this requirement, and transcendental outputs at exactly the bases that would meet it. The classification is therefore self-consistent with a counterexample rather than obstructive to one: it converts every attempt to manufacture a third independent base into the very transcendence statement one was trying to prove. The kernel-verified radical work makes the same wall concrete: the simultaneous normalization of the two canonical radicals produces an explicit multiplicative relation among 2, the primitive odd core, and the swapped radical, so that the most tempting three-base application of the algebraic six-exponentials theorem fails its monomial-injectivity hypothesis for a structural, kernel-verified reason (`swappedRadical_relation_blocks_monomial_injectivity` in `CanonicalRadicalFieldNorm`). Closing the gap needs a new source of size — nonarchimedean, or from a sharper five/six-exponentials interface — not a wider supply of algebraic bases. The conjecture remains open.

14 Rational and integral outputs inside the spectrum

The algebraic spectrum identifies the divisible group

$$\Gamma = \{2^r 3^s : r, s \in \mathbb{Q}\}$$

as the exact set of positive algebraic bases whose powers at a hypothetical nonintegral solution remain algebraic: outside Γ every such power is transcendental, and inside Γ every such power is algebraic. A finer arithmetic question survives inside the spectrum. If $\alpha = 2^r 3^s \in \Gamma$, when is α^x rational, and when is it a positive integer?

For a fixed solution the answer is completely encoded by prime valuations of the output pair. Rationality becomes a congruence condition, integrality adds a cone condition, and the gap between the obvious rational bases $2^u 3^v$ with $u, v \in \mathbb{Z}$ and the full rational locus is one explicit finite abelian group whose order is a single gcd of 2×2 minors. This is a strict refinement of the divisible-hull description of the algebraic locus, and it is exact rather than asymptotic.

14.1 The valuation matrix of an output pair

Let $x \in \mathcal{S} \setminus \mathbb{Z}$ be a nonintegral solution and write

$$M = 2^x \in \mathbb{N}, \quad A = 3^x \in \mathbb{N}.$$

Let P be the finite union of the prime supports of M and A , and form the column vectors

$$\mathbf{m} = (v_p(M))_{p \in P}, \quad \mathbf{a} = (v_p(A))_{p \in P}, \quad V_x = \begin{pmatrix} \mathbf{m} & \mathbf{a} \end{pmatrix} \in M_{P \times 2}(\mathbb{Z}).$$

For a column $q = (r, s)^T \in \mathbb{Q}^2$ put $\alpha(q) = 2^r 3^s$, so that Γ is exactly the image of $q \mapsto \alpha(q)$. Taking the unique positive real interpretation of the rational powers,

$$\alpha(q)^x = (2^x)^r (3^x)^s = M^r A^s = \prod_{p \in P} p^{rv_p(M) + sv_p(A)}. \quad (28)$$

The exponent vector appearing in (28) is precisely $V_x q$.

Lemma 14.1 (The valuation matrix has rank two). *The two columns of V_x are linearly independent over $\mathbb{Q}$.*

Proof. Suppose $r\mathbf{m} + s\mathbf{a} = 0$ with $r, s \in \mathbb{Q}$. Then every prime exponent in (28) vanishes, so $M^r A^s = 1$. Taking real logarithms and using $\log M = x \log 2$, $\log A = x \log 3$ gives $x(r \log 2 + s \log 3) = 0$. Since a solution with $M = 2^x \in \mathbb{N}$, $M > 1$ has $x > 0$, and since $\vartheta = \log_2 3$ is irrational, $r = s = 0$. $\square$

Proposition 14.2 (Exact rational and integral auxiliary-output spectrum). *Let $x \in \mathcal{S} \setminus \mathbb{Z}$, let V_x be as above, and let $q \in \mathbb{Q}^2$. Then*

$$\alpha(q)^x \in \mathbb{Q}_{>0} \iff V_x q \in \mathbb{Z}^P, \quad \alpha(q)^x \in \mathbb{N} \iff V_x q \in \mathbb{N}_0^P.$$

Together with the algebraic spectrum, this classifies algebraic, rational and integral outputs for every positive algebraic auxiliary base. [paper proof in this report]

Proof. Equation (28) exhibits $\alpha(q)^x$ as a finite product of primes with rational exponents. Choose a common denominator d for the entries of $V_x q$. If $\alpha(q)^x$ is rational, then raising to the d -th power gives a rational number whose ordinary p -adic valuation equals d times the p -th entry of $V_x q$; comparing valuations of numerator and denominator shows every entry of $V_x q$ is an integer. Conversely, integer exponents produce a positive rational number. Such a rational number is a positive integer exactly when all of its prime valuations are nonnegative, which is the second equivalence. $\square$

Three nested spectra therefore occur, all determined by V_x :

$$\begin{aligned} \mathcal{B}_{\text{alg}}(x) &= \{\alpha(q) : q \in \mathbb{Q}^2\} = \Gamma, & \Lambda_x &= \{q \in \mathbb{Q}^2 : V_x q \in \mathbb{Z}^P\}, \\ \mathcal{B}_{\text{rat}}(x) &= \{\alpha(q) : q \in \Lambda_x\}, & \Lambda_x^+ &= \{q \in \Lambda_x : V_x q \geq 0\}. \\ \mathcal{B}_{\text{int}}(x) &= \{\alpha(q) : q \in \Lambda_x^+\}, \end{aligned}$$

The outer spectrum does not depend on x ; the two inner ones may.

14.2 Smith normal form and the finite rational-output defect

By Lemma 14.1 the matrix V_x has rank two, so Λ_x is a rank-two lattice in $\mathbb{Q}^2$ containing $\mathbb{Z}^2$ with finite index. Write

$$\delta_1(V_x) = \gcd\{\text{all entries of } V_x\}, \quad \delta_2(V_x) = \gcd \left\{ \left| \det \begin{pmatrix} v_p(M) & v_p(A) \\ v_{p'}(M) & v_{p'}(A) \end{pmatrix} \right| : p, p' \in P \right\}$$

for the two determinantal divisors. The second is positive precisely because the rank is two.

Proposition 14.3 (Smith classification of the rational locus). *Let $d_1 \mid d_2$ be the Smith invariant factors of V_x , so that $d_1 = \delta_1(V_x)$ and $d_1 d_2 = \delta_2(V_x)$. After a unimodular change of coordinates in $\mathbb{Q}^2$,*

$$\Lambda_x = \frac{1}{d_1} \mathbb{Z} \oplus \frac{1}{d_2} \mathbb{Z},$$

and consequently

$$\Lambda_x / \mathbb{Z}^2 \cong \mathbb{Z} / d_1 \mathbb{Z} \oplus \mathbb{Z} / d_2 \mathbb{Z}, \quad |\Lambda_x / \mathbb{Z}^2| = \delta_2(V_x). \quad [\text{paper proof in this report}]$$

Proof. Choose unimodular matrices $U \in \mathrm{GL}_P(\mathbb{Z})$ and $W \in \mathrm{GL}_2(\mathbb{Z})$ with

$$UV_xW = \begin{pmatrix} d_1 & 0 \\ 0 & d_2 \\ 0 & 0 \\ \vdots & \vdots \end{pmatrix}.$$

Since U preserves $\mathbb{Z}^P$, the condition $V_xq \in \mathbb{Z}^P$ is equivalent to $UV_xq \in \mathbb{Z}^P$, and substituting $q = Wq'$ turns it into $d_1q'_1 \in \mathbb{Z}$ and $d_2q'_2 \in \mathbb{Z}$. Since W preserves $\mathbb{Z}^2$, it identifies $\Lambda_x/\mathbb{Z}^2$ with the quotient computed in the primed coordinates, which is $\mathbb{Z}/d_1\mathbb{Z} \oplus \mathbb{Z}/d_2\mathbb{Z}$. The identifications $d_1 = \delta_1(V_x)$ and $d_1d_2 = \delta_2(V_x)$ are the standard determinantal-divisor formulae for the Smith normal form. $\square$

The first invariant factor has a direct output interpretation: d_1 is the largest integer d such that M and A are both perfect d -th powers, because $d \mid \delta_1(V_x)$ says exactly that d divides every prime valuation of M and of A . For the canonical least nonintegral solution β , primitivity of the least output pair (Section 5) gives $d_1 = 1$, and then

$$\Lambda_\beta/\mathbb{Z}^2 \cong \mathbb{Z}/\delta_2(V_\beta)\mathbb{Z}.$$

Thus a single integer — the gcd of all 2×2 valuation minors of the least output pair — measures the entire finite defect between the rational β -outputs at algebraic 2,3-radical bases and the obvious rational 2,3-unit bases. We call $\delta_2(V_\beta)$ the *rational-output defect*. It is a finite, in principle computable, invariant of a hypothetical counterexample, and it equals 1 exactly when no auxiliary base outside $\{2^u3^v : u, v \in \mathbb{Z}\}$ produces a rational output.

The integral spectrum is the intersection of this lattice with the rational polyhedral cone $V_xq \geq 0$. By Gordan’s lemma Λ_x^+ is a finitely generated normal affine semigroup, so $\mathcal{B}_{\mathrm{int}}(x)$ has a finite generating set of auxiliary bases. This is an exact structural replacement for attempting to enumerate algebraic auxiliary bases one at a time.

Status: [\[kernel-verified Lean\]](#) and [\[paper proof in this report\]](#)

The lattice statements themselves — Lemma 14.1, Proposition 14.2, Proposition 14.3 and the Gordan consequence — are [\[paper proof in this report\]](#): they are proved on paper in this report and are not yet formalized in that form. What *is* kernel-verified is the arithmetic core on which they rest, in three modules.

`CanonicalRadicalValuation` proves the divisibility dictionary between perfect powers and valuations that underlies the meaning of d_1 : a divisor of the odd-prime valuation gcd of a positive integer is an honest exponent of an exact power

(`exists_pow_eq_of_dvd_oddPrimeValuationGCD`), the least rational-power index is coprime to the joint gcd of the denominator exponent and the numerator’s valuation gcd

(`rationalPowerIndex_gcd_denominator_valuationGCD_eq_one`), and every prime dividing that index leaves some prime valuation of the normalized numerator undivided, giving the exact radical degree together with the valuation restriction

(`exists_exact_radical_degree_and_valuation_restriction`).

`SimultaneousRadicalDegrees` kernel-verifies the exactness statements that the $d_1 = 1$ normalization buys: under the simultaneous primitive normalization the displayed exponents are the least rational-power indices and hence the exact degrees over $\mathbb{Q}$

(`oddCoreRpow_exact_rationalPowerIndex_of_simultaneous_normalization` and `threeFreeCoreRpow_exact_radical_degree_of_simultaneous_normalization`), and

when both base-adic depths vanish the two indices are coprime and the mixed product has exact degree equal to their product

(`mixed_coreRpow_exact_radical_degree_of_zero_baseDepths`). The descent used there is precisely the one behind $d_1 = 1$: a proper index divisor would make both outputs of the least nonintegral solution common perfect powers, which is excluded by

`IsLeastTwoBaseNonintegerSolution.no_common_perfect_power` in `LeastSolution` (Section 5). That exclusion is exactly the assertion $d_1 = 1$ for V_β .

`CanonicalRadicalMixedOutput` kernel-verifies the transcendence side of the mixed case: a genuinely mixed positive monomial in the two canonical radicals retains an external prime in every rational power, so all of its outputs at the least nonintegral solution are transcendental, and an algebraic mixed output is therefore forced onto a coordinate face. The theorem is

`TwoBaseNonintegerSolution.transcendental_mixed_canonical_radical_output`,

packaged under failure of the conjecture as

`exists_simultaneous_canonical_radical_mixed_output_transcendence`.

These are statements about the canonical radicals rather than about Λ_x itself; they confirm the geometry of the lattice picture without yet formalizing it.

Exact simultaneous radical degrees and mixed-output exclusion. Under failure of the conjecture, let β be the least nonintegral solution and choose its simultaneous primitive normalization

$$2^\beta = 2^a w^d, \quad 3^\beta = 3^b v^e, \quad a = 0 \text{ or } b = 0.$$

For

$$E = w^\vartheta, \quad F = v^{1/\vartheta},$$

the simultaneous power identities are

$$E^d = \frac{3^b v^e}{3^a}, \quad F^e = \frac{2^a w^d}{2^b}.$$

The least positive rational-power indices of E and F are exactly d and e . Consequently their minimal polynomials are the displayed pure binomials

$$X^d - \frac{3^b v^e}{3^a}, \quad X^e - \frac{2^a w^d}{2^b},$$

and their degrees over $\mathbb{Q}$ are exactly d and e . In the zero-depth case $a = b = 0$ one has $\gcd(d, e) = 1$, and for $G = EF$,

$$\text{minpoly}_{\mathbb{Q}}(G) = X^{de} - v^{e^2} w^{d^2}, \quad [\mathbb{Q}(G) : \mathbb{Q}] = de.$$

Moreover, the canonical data supplied by failure contain an external prime, and for every $i, j > 0$ the mixed output $(E^i F^j)^\beta$ is transcendental; if a nonnegative mixed output is algebraic, its exponent pair lies on one of the two coordinate axes. [\[kernel-verified Lean\]](#)

These conclusions are mutually compatible, not contradictory. The degree statements concern the algebraic numbers E , F , and EF , whereas the real irrational-power outputs at β are transcendental. A field embedding of $\mathbb{Q}(E, F)$ need not commute with the chosen real operation $z \mapsto z^\beta$, so taking conjugates or norms does not turn the mixed-output transcendence into a forbidden algebraic relation. In the zero-depth case the product degree de is the maximal degree allowed by the two coprime pure-radical degrees, not an impossible dependence. Thus the simultaneous theory sharpens the exact algebraic data and excludes the interior of the mixed-output quadrant, but leaves both coordinate axes — and the full Alaoglu–Erdős conjecture — open.

14.3 Two translated solutions remove the fractional defect

The algebraic spectrum is solution-independent, but the rational locus Λ_x need not be. Translation of the exponent by one gives a sharp intersection identity, and it shows that the defect cannot persist across the whole solution monoid.

Proposition 14.4 (Translation intersection). *Let $x \in \mathcal{S} \setminus \mathbb{Z}$. Then*

$$\Lambda_x \cap \Lambda_{x+1} = \mathbb{Z}^2.$$

Equivalently, for $\alpha \in \overline{\mathbb{Q}}_{>0}$, if α^x and α^{x+1} are both rational, then $\alpha = 2^u 3^v$ with $u, v \in \mathbb{Z}$. [\[paper proof in this report\]](#)

Proof. One inclusion is immediate: for $q \in \mathbb{Z}^2$ the base $\alpha(q) = 2^u 3^v$ has $\alpha(q)^y = M^u A^v \in \mathbb{Q}_{>0}$ for every solution y , in particular for $y = x$ and $y = x + 1$. Conversely, suppose α^x and α^{x+1} are rational. Both are algebraic, so the algebraic spectrum gives $\alpha = 2^r 3^s$ with $r, s \in \mathbb{Q}$, and

$$\alpha = \frac{\alpha^{x+1}}{\alpha^x} \in \mathbb{Q}_{>0}.$$

A positive rational of the form $2^r 3^s$ with $r, s \in \mathbb{Q}$ has $r, s \in \mathbb{Z}$, by comparing ordinary 2-adic and 3-adic valuations after clearing a common denominator. $\square$

Corollary 14.5 (Universal rational and integral spectra under failure). *Assume Conjecture 2.1 fails and let $\alpha \in \overline{\mathbb{Q}}_{>0}$. [paper proof in this report]*

1. α^y is rational for every nonintegral solution y if and only if $\alpha = 2^u 3^v$ with $u, v \in \mathbb{Z}$.
2. α^y is a positive integer for every solution y if and only if $\alpha = 2^u 3^v$ with $u, v \in \mathbb{N}_0$.

Proof. For the first part, apply Proposition 14.4 to a nonintegral solution x and to $x + 1$, which is again a nonintegral solution. Conversely, rational 2, 3-units have rational outputs $M^u A^v$ at every solution.

For the second part, assume $\alpha^y \in \mathbb{N}$ for every $y \in \mathcal{S}$. Taking $y = x$ and $y = x + 1$ makes $\alpha = \alpha^{x+1}/\alpha^x$ a positive rational, say a/b in lowest terms. Put $z = \alpha^x \in \mathbb{N}$. Every $x + n$ with $n \geq 0$ is again a solution, so the number $z(a/b)^n = \alpha^{x+n}$ is an integer for every n , whence b^n divides the fixed integer z for every n , forcing $b = 1$ and $\alpha \in \mathbb{N}$. Since α^x is an integer, the algebraic spectrum gives $\alpha \in \Gamma$, and an integer in Γ is exactly $2^u 3^v$ with $u, v \in \mathbb{N}_0$. The converse is immediate. $\square$

Corollary 14.5 says that the rational-output defect is a genuinely per-solution phenomenon. A base can enjoy a rational output at one nonintegral solution only by failing at the next one, unless it is a rational 2, 3-unit already.

14.4 The two smooth branches are mutually exclusive

The number of rows of V_x is another exact invariant of a counterexample, and the smallest conceivable case is excluded outright. Call a positive integer $\{2, 3\}$ -smooth (equivalently 3-smooth) when it is divisible by no prime outside $\{2, 3\}$. In the notation above, M is $\{2, 3\}$ -smooth exactly when the column $\mathbf{m}$ is supported on $\{2, 3\}$, and both outputs are $\{2, 3\}$ -smooth exactly when $P \subseteq \{2, 3\}$, that is, when V_x is a 2×2 integer matrix.

Proposition 14.6 (The smooth cases are mutually exclusive). *For a nonintegral solution x , the outputs $M = 2^x$ and $A = 3^x$ cannot both be $\{2, 3\}$ -smooth. Equivalently, the valuation matrix V_x of a counterexample always has a row indexed by a prime $p \geq 5$. [kernel-verified Lean]*

Proof. Suppose $M = 2^a 3^b$ and $A = 2^c 3^e$ with $a, b, c, e \in \mathbb{N}_0$. Taking base-two logarithms of $2^x = M$ and of $3^x = A$ gives the two relations

$$x = a + b\vartheta, \quad x\vartheta = c + e\vartheta, \quad \vartheta = \log_2 3.$$

Substituting the first into the second eliminates x and leaves

$$b\vartheta^2 + (a - e)\vartheta - c = 0.$$

Were any of b , $a - e$, c nonzero, this would exhibit ϑ as a root of a nonzero rational polynomial of degree at most two, contradicting its transcendence. Hence $b = c = 0$ and $a = e$, so $x = a \in \mathbb{Z}$, contrary to hypothesis. $\square$

Corollary 14.7 (A smooth output has a rough partner). *Let x be a nonintegral solution with $M = 2^x$ $\{2, 3\}$ -smooth. Then $A = 3^x$ has a prime divisor $p \geq 5$. [\[kernel-verified Lean\]](#)*

Proof. If every prime divisor of A were 2 or 3, then A would be $\{2, 3\}$ -smooth, and Proposition 14.6 would force $x \in \mathbb{Z}$. $\square$

Status: [\[kernel-verified Lean\]](#)

`integer_of_threeSmooth_outputs` and `exists_prime_five_le_dvd_of_threeSmooth_candidate` in `SmoothOutputs`, with the integral restatement `integer_of_two_three_rpow_integer_of_threeSmooth_outputs`. The proofs use only the transcendence of ϑ , itself kernel-verified in this corpus from Gelfond–Schneider. The same statement is also derived independently at paper level in the research report by the identical degree-drop elimination, so the formalized and unformalized routes confirm each other.

Two remarks fix the position of this exclusion.

The two-sided hypothesis is essential. Assuming only that M is $\{2, 3\}$ -smooth gives $x = a + b\vartheta$ and then requires $(3^\vartheta)^b$ to be rational up to a power of 3, which is exactly the open localized-radical condition of Section 6. What the second smoothness hypothesis supplies is a *second* linear relation between x and ϑ ; two such relations eliminate x and leave an algebraic equation in ϑ alone. This degree drop is the entire mechanism, and it is worth isolating as a template: each independent multiplicative constraint on an output buys one linear relation, and two relations force algebraicity of ϑ . The difficulty of the general problem is precisely that a second constraint is never free.

The exclusion feeds back into the lattice. Because $P \not\subseteq \{2, 3\}$, the valuation matrix of a counterexample is never square, so $\delta_2(V_x)$ is a genuine gcd over at least three rows rather than a single determinant, and the rational locus Λ_x is cut out by at least three congruences. In particular the rational-output defect is bounded by the absolute value of any one 2×2 minor of V_x , and any two minors that are coprime force $\delta_2(V_x) = 1$, in which case no auxiliary base outside $\{2^u 3^v : u, v \in \mathbb{Z}\}$ has a rational output at x at all.

15 The second-iterate kernel and multiplicative dependence

The algebraic spectrum of Section 13 controls one application of the exponent x . Asking which bases of Γ return to Γ after that application turns second iterates into finite-dimensional linear algebra, and — as this section records — into an entirely elementary statement about four integers.

15.1 The kernel and its invariance

For a nonintegral solution x with outputs $M = 2^x$ and $A = 3^x$, set

$$K_x = \{(r, s) \in \mathbb{Q}^2 : M^r A^s \in \Gamma\}, \quad \Gamma = \{2^u 3^v : u, v \in \mathbb{Q}\}. \quad (29)$$

Writing $\mathbf{m}_{\text{ext}} = (v_p(M))_{p \neq 2,3}$ and $\mathbf{a}_{\text{ext}} = (v_p(A))_{p \neq 2,3}$ for the external valuation vectors, membership of $M^r A^s$ in Γ is exactly the vanishing of every prime valuation outside $\{2, 3\}$, so

$$K_x = \ker((r, s) \mapsto r \mathbf{m}_{\text{ext}} + s \mathbf{a}_{\text{ext}}). \quad (30)$$

Proposition 15.1 (Kernel dimension). $\dim_{\mathbb{Q}} K_x \leq 1$; the kernel contains no vector with $rs > 0$; it equals $\mathbb{Q}(1, 0)$ exactly when M is 3-smooth and $\mathbb{Q}(0, 1)$ exactly when A is 3-smooth.

Proof. At least one external vector is nonzero: otherwise both M and A would be 3-smooth, which Theorem 12.5 excludes. Hence the map in (30) has rank at least one and its kernel has dimension at most one. If $r, s > 0$ then any external prime dividing M or A has positive valuation in $M^r A^s$, and if $r, s < 0$ it has negative valuation, so no same-sign vector lies in the kernel. The two smooth cases are immediate from (30) by testing $(1, 0)$ and $(0, 1)$. $\square$

Status: [\[kernel-verified Lean\]](#) — `SecondIterateKernel` carries the kernel as an integral subgroup `kernelAddSubgroup` of $\mathbb{Z}^2$, with `kernelPairs_mul_nonpos` for the absence of same-sign vectors, `kernelPairs_det_eq_zero` and `kernelPairs_prop_dim` for the rank bound, `mem_kernelPairs_one_zero_iff` for the two smooth cases, and `saturatedKernel_invariant` for $K_x = K_y$. The input that M and A cannot both be 3-smooth is `integer_of_threeSmooth_outputs` in `SmoothOutputs`. Clearing denominators lets the formal statement use integer rather than rational pairs, which is equivalent and lighter.

The kernel is not an artifact of the chosen counterexample. If x, y are two nonintegral solutions and $\alpha = 2^r 3^s$, then $B = \alpha^x$ is algebraic, and applying the spectrum theorem to y gives $\alpha^{xy} = B^y \in \overline{\mathbb{Q}}$ if and only if $B \in \Gamma$, that is, if and only if $(r, s) \in K_x$. The left-hand condition is symmetric in x and y , so $K_x = K_y$: the kernel is an invariant of the problem, not of the solution.

15.2 Multiplicative dependence of four integers

The kernel condition has a formulation with no logarithms, no projective linear group, and no transcendence vocabulary at all. It is the observation of this section.

Theorem 15.2 (Multiplicative form of the kernel dichotomy). *Let x be a nonintegral solution with outputs $M = 2^x$ and $A = 3^x$. The following are equivalent:*

- (i) $K_x \neq \{0\}$;
- (ii) the four integers $M, A, 2, 3$ are multiplicatively dependent, that is,

$$M^a A^b 2^c 3^d = 1 \quad \text{for some } (a, b, c, d) \in \mathbb{Z}^4 \setminus \{0\};$$

- (iii) $1, \vartheta, x, x\vartheta$ are linearly dependent over $\mathbb{Q}$;
- (iv) x lies in the $\text{PGL}_2(\mathbb{Q})$ -orbit of ϑ , that is, $x = (u + v\vartheta)/(r + s\vartheta)$ with rational u, v, r, s and $us - vr \neq 0$.

Proof. (i) implies (ii). If $(r, s) \in K_x$ is nonzero then $M^r A^s = 2^u 3^v$ with $u, v \in \mathbb{Q}$; clearing a common denominator of r, s, u, v produces integers with $M^a A^b 2^c 3^d = 1$ and $(a, b) \neq (0, 0)$.

(ii) implies (i). Given such a relation, note first that $(a, b) \neq (0, 0)$: otherwise $2^c 3^d = 1$, and unique factorization forces $c = d = 0$, contradicting nontriviality. Then $M^a A^b = 2^{-c} 3^{-d} \in \Gamma$, so $(a, b) \in K_x \setminus \{0\}$.

(ii) is equivalent to (iii). Taking logarithms of $M^a A^b 2^c 3^d = 1$ and using $\log M = x \log 2$ and $\log A = x \log 3$ gives

$$(ax + c) \log 2 + (bx + d) \log 3 = 0,$$

and division by $\log 2$ turns this into

$$x(a + b\vartheta) + (c + d\vartheta) = 0, \tag{31}$$

which is precisely a nontrivial rational linear relation among $1, \vartheta, x, x\vartheta$. Each step reverses.

(iii) is equivalent to (iv). In (31) the coefficient $a + b\vartheta$ cannot vanish, since ϑ is irrational and $(a, b) \neq (0, 0)$; hence $x = -(c + d\vartheta)/(a + b\vartheta)$. If the determinant $ad - bc$ vanished, numerator and denominator would be rationally proportional and x would be rational, hence integral by Lemma 2.3, contrary to hypothesis. Conversely, any such expression cross-multiplies to (31). $\square$

Status: [\[kernel-verified Lean\]](#) — KernelDichotomy formalizes the equivalence of (ii), (iii) and (iv), as `multiplicativelyDependentOutputs_iff_secondIterateKernelRelation` together with `exists_moebius_of_multiplicativelyDependentOutputs` and `multiplicativelyDependentOutputs_of_moebius`; the packaged branching is `twoBaseNonintegerSolution_kernel_dichotomy`. Clause (i) is `SecondIterateKernel`. The formalization is slightly stronger than the statement above: (ii) $\iff$ (iii) needs only irrationality of ϑ , and nonintegrality of x enters just for the nonvanishing of the Möbius determinant.

15.3 Why the multiplicative form is worth having

Two consequences follow from stating the dichotomy multiplicatively.

It removes the transcendence vocabulary from the branching. Every counterexample falls into exactly one of two cases, and the case distinction is a question about four ordinary integers: either they satisfy a multiplicative relation or they do not. In the dependent branch Proposition 15.1 says the relation is essentially unique up to scaling, and its exponents on M and A have opposite signs unless one of the outputs is 3-smooth.

It makes Baker’s theorem applicable to the independent branch. This is the point that the Diophantine formulations obscure. In the independent branch $M, A, 2, 3$ are multiplicatively independent, so $\log M, \log A, \log 2, \log 3$ are $\mathbb{Q}$ -linearly independent *logarithms of algebraic numbers* — all four arguments are positive integers. Baker’s homogeneous theorem then gives more than $\mathbb{Q}$ -linear independence: these four logarithms are linearly independent over $\overline{\mathbb{Q}}$. The corpus has not previously used this input.

Corollary 15.3 (Algebraic-coefficient rigidity in the independent branch). *Suppose x is a nonintegral solution for which $M, A, 2, 3$ are multiplicatively independent. Then no relation*

$$\gamma_1 \log M + \gamma_2 \log A + \gamma_3 \log 2 + \gamma_4 \log 3 = 0$$

holds with algebraic $\gamma_1, \gamma_2, \gamma_3, \gamma_4$ not all zero.

Status: [classical/open background] — immediate from Baker’s homogeneous theorem on $\overline{\mathbb{Q}}$ -linear independence of logarithms of algebraic numbers, applied to the four positive integers $M, A, 2, 3$.

It must be said plainly what this does and does not do. It does *not* close the independent branch: the defining relation $\log M \log 3 = \log A \log 2$ has coefficients $\log 3$ and $-\log 2$, which are transcendental by Lindemann, so it is not a $\overline{\mathbb{Q}}$ -linear relation and Corollary 15.3 does not apply to it. That is exactly the “Baker is linear, not bilinear” obstruction recorded in Appendix L. What changes is that the independent branch is no longer structureless: it carries a genuine unconditional rigidity statement, and any future argument that manages to produce an algebraic-coefficient relation among these four logarithms — for instance by clearing the transcendental coefficients against a second, independent relation — would close it immediately.

16 Depth-three tower rigidity

A hypothetical counterexample makes the first level of the tower $2 \mapsto 2^x$ integral by assumption. This section answers the exact question that the corpus keeps circling: starting from an arbitrary positive algebraic base, how many times can one apply the same nonintegral exponent x before transcendence is forced? The answer is at most three, the bound is attained, and the depth at which the first transcendental value appears is determined by a single rational-linear invariant.

For positive real bases there is no branch ambiguity between iterated powers and powers of powers, so the tower is unambiguous:

$$((\gamma^x)^x)^x = \gamma^{x^3}.$$

Throughout this section x denotes a hypothetical nonintegral solution, that is $x \in \mathcal{S} \setminus \mathbb{N}_0$, and we write

$$m = 2^x \in \mathbb{N}, \quad n = 3^x \in \mathbb{N}.$$

The letter γ is reserved for a positive algebraic base, matching the Lean naming in `AlgebraicBaseUnitPower`.

16.1 The saturated spectrum and its strong form

Give $\overline{\mathbb{Q}}_{>0}$ its structure of $\mathbb{Q}$ -vector space, with rational scalars acting through the unique positive real root, $q \cdot \gamma = \gamma^q$, and set

$$\Gamma = \{2^r 3^s : r, s \in \mathbb{Q}\} \subset \mathbb{R}_{>0}.$$

This is the divisible hull of the group generated by 2 and 3: a countable, divisible subgroup of rational rank two, dense in $\mathbb{R}_{>0}$ already through $\{2^r : r \in \mathbb{Q}\}$, and isomorphic to $\mathbb{Q}^2$ as a $\mathbb{Q}$ -vector space via $(r, s) \mapsto 2^r 3^s$. An *integer* lies in Γ exactly when it is three-smooth, since clearing denominators in $N = 2^r 3^s$ gives $N^k = 2^p 3^q$ with $p, q \in \mathbb{Z}$ and unique factorization then confines the primes of N to $\{2, 3\}$.

Theorem 16.1 (Exact algebraic-base spectrum). *Let $x \in \mathcal{S} \setminus \mathbb{N}_0$ and let $\gamma \in \overline{\mathbb{Q}}_{>0}$. Then*

$$\gamma^x \in \overline{\mathbb{Q}} \iff \gamma \in \Gamma.$$

In particular the algebraic spectrum of a counterexample does not depend on which nonintegral solution is used.

Status: [\[kernel-verified Lean\]](#) — in `AlgebraicBaseUnitPower`, the pair

`TwoBaseNonintegerSolution.real_rpow_isAlgebraic_iff_hasPositiveTwoThreeUnitPower`

and

`TwoBaseNonintegerSolution.transcendental_real_rpow_of_no_positiveTwoThreeUnitPower`

The Lean predicate is `HasPositiveTwoThreeUnitPower`: some positive natural power of γ is a positive rational 2, 3-unit. For positive algebraic γ this is exactly membership in Γ — one direction raises $2^r 3^s$ to a common denominator, the other takes the unique positive real root — so the kernel-verified biconditional is the displayed spectrum verbatim.

The spectrum is already a strong statement: two algebraic outputs exhaust the maximal algebraic multiplicative rank that the two-base configuration can support, and every positive algebraic base outside one explicit dense subgroup produces a transcendental value. For the tower we need the sharper conclusion that the offending *logarithm* is not merely irrational but escapes the algebraic span of all logarithms.

Let $\mathcal{L}$ denote the $\overline{\mathbb{Q}}$ -vector space spanned by 1 and by all logarithms of nonzero algebraic numbers, arbitrary complex branches allowed; for a positive real algebraic number the real logarithm is one of those branches.

Definition 16.2. A positive real value e^z is *log-strongly transcendental* when $z \notin \mathcal{L}$. This implies ordinary transcendence, because the logarithm of a positive algebraic number lies in $\mathcal{L}$ by definition.

Theorem 16.3 (Strong algebraic-base spectrum). *Let $x \in \mathcal{S} \setminus \mathbb{N}_0$ and let $\gamma \in \overline{\mathbb{Q}}_{>0} \setminus \Gamma$. Then*

$$x \log \gamma \notin \mathcal{L},$$

so γ^x is log-strongly transcendental. The same conclusion holds with 2, 3 replaced by any multiplicatively independent pair $a, b \in \overline{\mathbb{Q}}_{>0}$ with $a^x, b^x \in \mathbb{Q}$, and Γ replaced by $\Gamma(a, b) = \{a^r b^s : r, s \in \mathbb{Q}\}$.

Proof. The logarithms $\log 2, \log 3, \log \gamma$ are linearly independent over $\mathbb{Q}$: a rational relation clears to an integral multiplicative relation, which either contradicts the multiplicative independence of 2 and 3 or places γ in Γ . Baker’s homogeneous theorem [3] upgrades this to linear independence over $\overline{\mathbb{Q}}$.

The exponent x is transcendental: it is irrational by the kernel-verified rational-exponent lemma of Section 4, and an algebraic irrational x would make 2^x transcendental by Gelfond–Schneider [9]. Hence 1 and x are linearly independent over $\overline{\mathbb{Q}}$.

Apply Roy’s strong six exponentials theorem [8, 11] to the three rows $\log 2, \log 3, \log \gamma$ and the two columns 1, x . Five of the six pairwise products lie in $\mathcal{L}$, namely

$$\log 2, \quad \log 3, \quad \log \gamma, \quad x \log 2 = \log m, \quad x \log 3 = \log n.$$

Therefore the sixth product $x \log \gamma$ lies outside $\mathcal{L}$. □

Status: [\[paper proof in this report\]](#) — paper proof in this report. It rests on Baker’s theorem and on Roy’s strong form of the six exponentials theorem, neither of which is formalized in the repository; the repository’s own interpolation determinant is specialized to integer entries. Theorem 16.1, the ordinary-transcendence shadow of this statement, is [\[kernel-verified Lean\]](#).

16.2 The return kernel

The spectrum controls one application of x . The tower needs to know which bases *inside* Γ come back to Γ after that application. On prime valuations this is finite-dimensional linear algebra.

Proposition 16.4 (The two outputs are not both three-smooth). *For $x \in \mathcal{S} \setminus \mathbb{N}_0$, at most one of $m = 2^x$ and $n = 3^x$ is three-smooth.*

Proof. Suppose $m = 2^a 3^b$ and $n = 2^c 3^d$ with $a, b, c, d \in \mathbb{N}_0$. Taking logarithms in base 2 and in base 3 gives $x = a + b\vartheta$ and $x = d + c/\vartheta$; eliminating x and clearing the denominator yields

$$b\vartheta^2 + (a - d)\vartheta - c = 0. \quad (32)$$

The kernel-verified transcendence of ϑ forces $b = c = 0$ and $a = d$, hence $x = a \in \mathbb{N}_0$, contrary to hypothesis. $\square$

Status: [\[paper proof in this report\]](#) — the only nonelementary input is the kernel-verified transcendence of ϑ ,

`transcendental_logThreeDivLogTwo`

in **Transcendence**. This is the same computation as the three-smooth output test of Section 9, restated here in the form the kernel below consumes.

For $x \in \mathcal{S} \setminus \mathbb{N}_0$ define the *return kernel*

$$K_x = \{\gamma \in \Gamma : \gamma^x \in \Gamma\} = \{2^r 3^s : (r, s) \in \mathbb{Q}^2, m^r n^s \in \Gamma\}. \quad (33)$$

Under the coordinate isomorphism $\Gamma \cong \mathbb{Q}^2$ this is a $\mathbb{Q}$ -subspace, because $(2^r 3^s)^x = m^r n^s$ is multiplicative in the coordinates. Write

$$\mathbf{m}_{\text{ext}} = (v_p(m))_{p \neq 2, 3}, \quad \mathbf{n}_{\text{ext}} = (v_p(n))_{p \neq 2, 3},$$

two finitely supported nonnegative integer vectors.

Proposition 16.5 (Kernel formula and dimension). *In the coordinates (r, s) ,*

$$K_x = \ker\left(\mathbb{Q}^2 \longrightarrow \mathbb{Q}^{\mathbb{P} \setminus \{2, 3\}}, \quad (r, s) \longmapsto r \mathbf{m}_{\text{ext}} + s \mathbf{n}_{\text{ext}}\right). \quad (34)$$

Moreover:

- (i) $\dim_{\mathbb{Q}} K_x \leq 1$;

- (ii) K_x contains no $2^r 3^s$ with $rs > 0$;
- (iii) $K_x = \{2^r : r \in \mathbb{Q}\}$ exactly when m is three-smooth;
- (iv) $K_x = \{3^s : s \in \mathbb{Q}\}$ exactly when n is three-smooth.

Proof. Membership of $m^r n^s$ in Γ is precisely the vanishing of every prime valuation outside $\{2, 3\}$, which is (34). By Proposition 16.4 at least one of the two external vectors is nonzero, so the displayed linear map has rank at least one and its kernel has dimension at most one.

If $r, s > 0$, then any external prime p occurring in m or in n has $rv_p(m) + sv_p(n) > 0$, since both valuations are nonnegative and at least one is positive; the case $r, s < 0$ is symmetric. Hence no same-sign direction survives.

Finally, if m is three-smooth then $\mathbf{m}_{\text{ext}} = 0$ while $\mathbf{n}_{\text{ext}} \neq 0$, so the kernel is exactly the first coordinate axis; the converse follows by testing the direction $(1, 0)$, which lies in K_x exactly when $m \in \Gamma$, that is, exactly when the integer m is three-smooth. The statement for n is the mirror image. $\square$

Status: [paper proof in this report] — elementary given Proposition 16.4. No transcendence input beyond the transcendence of ϑ enters the proof; the six exponentials theorem is used only to interpret K_x , not to compute it.

The dynamical content of the kernel is that it admits no nontrivial internal orbit of length two.

Lemma 16.6 (No two-step internal orbit). *Let $x \in \mathcal{S} \setminus \mathbb{N}_0$ and let $\gamma \in K_x \setminus \{1\}$. Then $\gamma^x \notin K_x$; equivalently,*

$$\gamma \in \Gamma \text{ and } \gamma^x \in \Gamma \implies \gamma^{x^2} \notin \Gamma.$$

Proof. Suppose γ and $\gamma_1 = \gamma^x$ both lie in $K_x \setminus \{1\}$. By Proposition 16.5(i) the space K_x is at most one-dimensional over $\mathbb{Q}$, so two of its nonzero elements are rational scalar multiples of one another in the additive notation: $\gamma_1 = \gamma^q$ for some $q \in \mathbb{Q}$. Since also $\gamma_1 = \gamma^x$ and $\gamma \neq 1$, comparing logarithms gives $x = q$, contradicting the irrationality of x . Hence $\gamma_1 \notin K_x$, which by (33) says exactly that $\gamma_1^x = \gamma^{x^2} \notin \Gamma$. $\square$

There are therefore no cycles, no fixed rays, and no length-two paths lying entirely inside Γ : a rational direction may survive one return to the algebraic spectrum, but the returned direction cannot itself survive.

16.3 Exact survival depth

Theorem 16.7 (Depth-three algebraic-tower theorem). *Let $x \in \mathcal{S} \setminus \mathbb{N}_0$ and let $\gamma \in \overline{\mathbb{Q}}_{>0} \setminus \{1\}$. Then at least one of*

$$\gamma^x, \quad \gamma^{x^2}, \quad \gamma^{x^3}$$

is transcendental. More precisely, the first transcendental term occurs at exactly one of the following three depths.

- (1) **Depth one.** If $\gamma \notin \Gamma$, then γ^x is log-strongly transcendental.

(2) **Depth two.** If $\gamma \in \Gamma \setminus K_x$, then $\gamma^x \in \overline{\mathbb{Q}} \setminus \Gamma$ and γ^{x^2} is log-strongly transcendental.

(3) **Depth three.** If $\gamma \in K_x \setminus \{1\}$, then $\gamma^x \in \Gamma$, $\gamma^{x^2} \in \overline{\mathbb{Q}} \setminus \Gamma$, and γ^{x^3} is log-strongly transcendental.

Equivalently, for some $j \in \{1, 2, 3\}$ one has $x^j \log \gamma \notin \mathcal{L}$.

Proof. Depth one is Theorem 16.3.

Assume $\gamma \in \Gamma$, so that $\gamma_1 = \gamma^x$ is algebraic by Theorem 16.1. If $\gamma \notin K_x$, then $\gamma_1 \notin \Gamma$ by (33), and γ_1 is a positive algebraic base outside Γ ; the strong spectrum applied to γ_1 and the same exponent x makes $\gamma_1^x = \gamma^{x^2}$ log-strongly transcendental.

If instead $\gamma \in K_x \setminus \{1\}$, then $\gamma_1 \in \Gamma$, so $\gamma_2 = \gamma_1^x = \gamma^{x^2}$ is algebraic. By Lemma 16.6 we have $\gamma_1 \notin K_x$, hence $\gamma_2 \notin \Gamma$. One last application of the strong spectrum, now to the base γ_2 , gives $\gamma_2^x = \gamma^{x^3}$ log-strongly transcendental.

The three cases are mutually exclusive and exhaustive, and in each case all earlier terms of the tower are algebraic, so the stated depth is exact rather than an upper bound. $\square$

The classification is compact enough to tabulate. For a nonintegral solution x and a positive algebraic base $\gamma \neq 1$, the first level of the tower whose logarithm escapes $\mathcal{L}$ is:

Location of γ	First log outside $\mathcal{L}$	Structural reason
$\gamma \notin \Gamma$	$x \log \gamma$	exact algebraic spectrum
$\gamma \in \Gamma \setminus K_x$	$x^2 \log \gamma$	the first image leaves Γ
$\gamma \in K_x \setminus \{1\}$	$x^3 \log \gamma$	no two-step internal orbit

Status: [\[paper proof in this report\]](#) — paper proof in this report, resting on Theorem 16.3 and hence on Baker’s theorem and Roy’s strong six exponentials theorem. The ordinary-transcendence version of depth one is [\[kernel-verified Lean\]](#); the ordinary-transcendence version of depth two for the mixed integer bases $2^i 3^j$ with $i, j > 0$ is also [\[kernel-verified Lean\]](#), as recorded below.

The theorem places a strict ceiling on how far an algebraic power tower can imitate the integral first level of a counterexample. A counterexample buys two algebraic outputs; it never buys a third independent one, and the tower makes the deficit visible after at most three steps regardless of the base chosen.

16.4 Exact behaviour of the bases two, three, and six

Corollary 16.8 (Base-two tower dichotomy). *Let $x \in \mathcal{S} \setminus \mathbb{N}_0$ and $m = 2^x$.*

(a) *If m is not three-smooth, then 2^{x^2} is log-strongly transcendental.*

(b) *If m is three-smooth, say $m = 2^a 3^b$ with $a, b \in \mathbb{N}_0$, then*

$$2^{x^2} = m^x = (2^x)^a (3^x)^b = m^a n^b \in \mathbb{N}, \quad (35)$$

and 2^{x^3} is log-strongly transcendental.

The mirror statement holds with 2 replaced by 3 and m by n .

Proof. The base 2 lies in Γ , and its coordinate vector is $(1, 0)$, which lies in K_x exactly when $m \in \Gamma$, that is, by Proposition 16.5(iii), exactly when the integer m is three-smooth. Now apply Theorem 16.7: the non-smooth case is depth two and the smooth case is depth three. In the smooth case the intermediate value is computed by (35), where the middle equality is the identity $(2^a 3^b)^x = (2^x)^a (3^x)^b$ for positive real bases and the last membership holds because $m, n \in \mathbb{N}$. $\square$

Equation (35) is worth pausing on. In the smooth branch the second level of the tower is not merely algebraic but an explicit positive integer, expressed in the two original outputs with the same exponents that describe the first output. The tower does not become transcendental because the arithmetic runs out immediately; it becomes transcendental because the third level has nowhere left to land.

Corollary 16.9 (Base six always fails at depth two). *For every $x \in \mathcal{S} \setminus \mathbb{N}_0$ and all $r, s \in \mathbb{Q}$ of the same nonzero sign,*

$$x^2(r \log 2 + s \log 3) \notin \mathcal{L}.$$

In particular $x^2 \log 6 \notin \mathcal{L}$, so 6^{x^2} is log-strongly transcendental.

Proof. The base $\gamma = 2^r 3^s$ lies in Γ , and its coordinate vector (r, s) has positive coordinate product, so $\gamma \notin K_x$ by Proposition 16.5(ii). Its exact survival depth is therefore two, by Theorem 16.7(2). The case $r = s = 1$ is $\gamma = 6$. $\square$

This recovers and strengthens a kernel-verified family. The Lean theorem

`TwoBaseNonintegerSolution.transcendental_mixed_two_three_squared_exponents`

in `AlgebraicOutputLocus` proves that $(2^i 3^j)^{x^2}$ is transcendental for all natural $i, j > 0$, which is exactly the statement that no genuinely mixed three-smooth integer base survives to the second level. Corollary 16.9 extends that to all rational same-sign directions, including negative ones, and upgrades the conclusion from transcendence to escape from $\mathcal{L}$; more importantly it explains *why* the mixed bases fail, namely that the exceptional directions of K_x are confined to the two coordinate axes, which the mixed monomials avoid by construction. The kernel-verified case is the positive-integer corner of a statement that is really about a line in $\mathbb{Q}^2$.

Status: [\[paper proof in this report\]](#) for Corollaries 16.8 and 16.9 as stated.

[\[kernel-verified Lean\]](#) for the ordinary transcendence of $(2^i 3^j)^{x^2}$ with integer $i, j > 0$, in `AlgebraicOutputLocus`. The gap between the two is exactly the general six exponentials input discussed in Section 9.

16.5 One-base tower equivalences

The depth theorem converts into a reformulation of the conjecture that is striking in its indifference to the base.

Universal three-level detector (Theorem 16.10). Fix *any* positive algebraic $\gamma \neq 1$ — $\gamma = 5$, $\gamma = \sqrt{2}$, $\gamma = 2$, $\gamma = 3$, $\gamma = \frac{1}{2}(7 + 3\sqrt{5})$, anything at all. Conjecture 2.1 is equivalent to

$$\forall x \in \mathbb{R}, \quad 2^x, 3^x \in \mathbb{Z} \implies \gamma^x, \gamma^{x^2}, \gamma^{x^3} \in \overline{\mathbb{Q}},$$

and equally to the assertion that all three numbers $x \log \gamma$, $x^2 \log \gamma$, $x^3 \log \gamma$ lie in $\mathcal{L}$. No arithmetic feature of the base is used: three levels of the tower detect a counterexample for every positive algebraic base simultaneously. [\[paper proof in this report\]](#)

Theorem 16.10 (Universal three-level detector). *Fix $\gamma \in \overline{\mathbb{Q}}_{>0} \setminus \{1\}$. Then Conjecture 2.1 is equivalent to*

$$\forall x \in \mathbb{R}, \quad 2^x, 3^x \in \mathbb{Z} \implies \gamma^x, \gamma^{x^2}, \gamma^{x^3} \in \overline{\mathbb{Q}}, \quad (36)$$

and also to the same implication with $\overline{\mathbb{Q}}$ -algebraicity replaced by the three membership conditions $x^j \log \gamma \in \mathcal{L}$, $j = 1, 2, 3$.

Proof. If the conjecture holds, then $\mathcal{S} = \mathbb{N}_0$, every exponent x, x^2, x^3 appearing in (36) is a nonnegative integer, and all three powers of the algebraic number γ are algebraic, with logarithms in $\mathcal{L}$. If the conjecture fails, apply Theorem 16.7 to a nonintegral solution: one of the three values is log-strongly transcendental, so (36) fails and so does its $\mathcal{L}$ -version. $\square$

Bases already inside Γ have their first level algebraic for free, which removes one of the three conditions and yields the shortest form of the equivalence.

Corollary 16.11 (Base-two cubic-tower equivalence). *Conjecture 2.1 is equivalent to*

$$\forall x \in \mathbb{R}, \quad 2^x, 3^x \in \mathbb{Z} \implies 2^{x^2} \in \overline{\mathbb{Q}} \text{ and } 2^{x^3} \in \overline{\mathbb{Q}}. \quad (37)$$

The same statement holds with the base 3 in place of the base 2.

Proof. Take $\gamma = 2$ in Theorem 16.10 and drop the first condition, which is automatic because 2^x is an integer under the hypothesis. $\square$

This is a genuinely different equivalence from the second-level algebraicity tests of Section 9, which require *both* 2^{x^2} and 3^{x^2} to be algebraic. One original base suffices if one is willing to climb a single extra level. The trade is informative: a second base and a third tower level carry exactly the same amount of information, which is another way of seeing that the two-base configuration supplies only two independent logarithmic directions no matter how the exponent is iterated.

Status: [\[paper proof in this report\]](#) — paper proof in this report. Both equivalences inherit the classical inputs of Theorem 16.3; the ordinary-algebraicity forms (36) and (37) need only the six exponentials theorem, since Theorem 16.1 is [\[kernel-verified Lean\]](#) and only the depth-two and depth-three steps require the general form. Neither is a proof of Conjecture 2.1, which remains open: the tower reformulations relocate the difficulty to the second and third levels without supplying a mechanism that produces algebraicity there.

17 Symmetric tower constants

The exponent $\vartheta = \log_2 3$ and its reciprocal

$$\eta = \frac{1}{\vartheta} = \frac{\log 2}{\log 3}$$

generate two symmetric “next outputs”,

$$E_3 = 3^\vartheta = 5.704522494691117635\dots, \quad E_2 = 2^\eta = 1.548562652630242907\dots$$

The first is the canonical radical of Section 6 in the special branch where the primitive odd core equals 3; the second is its base-swapped mirror. Neither constant is known to be transcendental, and for E_3 even irrationality is open. What is available unconditionally is a statement about the *pair*, and what a counterexample would add is an exact classification of which of the two becomes algebraic.

17.1 An unconditional strong dichotomy

Theorem 17.1 (At least one symmetric tower constant is strongly transcendental). *At least one of the two numbers*

$$\vartheta \log 3 = \log E_3, \quad \eta \log 2 = \log E_2$$

lies outside $\mathcal{L}$. In particular, at least one of E_3 and E_2 is transcendental.

Proof. The three numbers $1, \vartheta, \eta$ are linearly independent over $\overline{\mathbb{Q}}$: a relation $a + b\vartheta + c\eta = 0$ with algebraic coefficients, multiplied by ϑ , becomes $b\vartheta^2 + a\vartheta + c = 0$, which is impossible for a transcendental ϑ unless $a = b = c = 0$. The two numbers $\log 2, \log 3$ are linearly independent over $\mathbb{Q}$ by unique factorization, hence over $\overline{\mathbb{Q}}$ by Baker’s theorem [3].

Apply Roy’s strong six exponentials theorem [8, 11] to the three rows $1, \vartheta, \eta$ and the two columns $\log 2, \log 3$. Four of the six pairwise products are already logarithms of algebraic numbers,

$$\log 2, \quad \log 3, \quad \vartheta \log 2 = \log 3, \quad \eta \log 3 = \log 2,$$

and therefore lie in $\mathcal{L}$. Hence at least one of the two remaining products, $\vartheta \log 3$ and $\eta \log 2$, lies outside $\mathcal{L}$. $\square$

Status: **[paper proof in this report]** — paper proof in this report, using Baker’s theorem and Roy’s strong six exponentials theorem. The ordinary six exponentials theorem gives the weaker dichotomy that one of E_3, E_2 is transcendental. The strong form is worth the extra input because it survives verbatim into the conditional branches below.

Note what the theorem does not say. It does not identify which constant is transcendental, and no known method identifies either one individually; that is precisely the wall described in Section 6. The value of the dichotomy is that a counterexample would resolve it, and would resolve it in a completely explicit way.

17.2 Exact conditional classification

Throughout this subsection $x \in \mathcal{S} \setminus \mathbb{N}_0$, $m = 2^x$, and $n = 3^x$.

Theorem 17.2 (The E_3 equivalences). *The following are equivalent.*

- (i) $E_3 = 3^\vartheta$ is algebraic;
- (ii) $\vartheta \log 3 \in \mathcal{L}$;
- (iii) m is three-smooth;
- (iv) there are $a, b \in \mathbb{N}_0$ with $b > 0$ such that $m = 2^a 3^b$ and $x = a + b\vartheta$;
- (v) $x \in \mathbb{Q} + \mathbb{Q}\vartheta$.

If these conditions fail, then $\vartheta \log 3 \notin \mathcal{L}$ and E_3 is log-strongly transcendental. If they hold, then

$$E_3^b = \frac{n}{3^a} \in \mathbb{Q}_{>0}, \quad (38)$$

so E_3 is an explicit algebraic radical.

Proof. (iv) $\Rightarrow$ (iii) $\Rightarrow$ (v): the first is trivial, and taking logarithms of $m = 2^a 3^b$ gives $x = a + b\vartheta$, which lies in $\mathbb{Q} + \mathbb{Q}\vartheta$.

(v) $\Rightarrow$ (iv): write $x = u + v\vartheta$ with $u, v \in \mathbb{Q}$ and clear a common positive denominator d , so that $m^d = 2^p 3^q$ with $p, q \in \mathbb{Z}$. The left side is a positive integer, so prime valuations force $p, q \geq 0$ and $d \mid p$, $d \mid q$. Hence $u = a$ and $v = b$ with $a, b \in \mathbb{N}_0$, and $b > 0$ because $x \notin \mathbb{N}_0$.

(i) $\Rightarrow$ (iv): apply the six exponentials theorem to the rows $1, x, \vartheta$ and the columns $\log 2, \log 3$. The six exponential values are

$$2, \quad 3, \quad m, \quad n, \quad 3, \quad E_3,$$

all algebraic by hypothesis. The theorem therefore forbids both independence conditions simultaneously; the columns $\log 2, \log 3$ are $\mathbb{Q}$ -independent, so the rows are $\mathbb{Q}$ -dependent and $x = u + v\vartheta$ for some $u, v \in \mathbb{Q}$. Now (v) $\Rightarrow$ (iv) applies.

(iv) $\Rightarrow$ (i): from $x = a + b\vartheta$ we get $n = 3^x = 3^a E_3^b$, which is (38) and exhibits E_3 as a positive real radical of a positive rational, hence algebraic.

Finally the strong statement. Suppose m is not three-smooth. Then $\log 2, \log m, \log 3$ are linearly independent over $\mathbb{Q}$: a rational relation with nonzero coefficient on $\log m$ would place m in Γ , and an integer in Γ is three-smooth, while a relation with zero coefficient on $\log m$ contradicts the independence of $\log 2, \log 3$. Baker's theorem upgrades this to independence over $\overline{\mathbb{Q}}$, so after division by $\log 2$ the three numbers $1, x, \vartheta$ are $\overline{\mathbb{Q}}$ -linearly independent. Apply Roy's theorem to the rows $1, x, \vartheta$ and the columns $\log 2, \log 3$: five of the six products, namely $\log 2, \log 3, \log m, \log n$, and $\vartheta \log 2 = \log 3$, lie in $\mathcal{L}$, so $\vartheta \log 3 \notin \mathcal{L}$. Thus failure of (iii) implies failure of (ii); and (i) implies (ii) because $\log E_3 = \vartheta \log 3$ is then the logarithm of an algebraic number. $\square$

Theorem 17.3 (The E_2 equivalences). *The following are equivalent.*

- (i) $E_2 = 2^\eta$ is algebraic;
- (ii) $\eta \log 2 \in \mathcal{L}$;

- (iii) n is three-smooth;
- (iv) there are $c, d \in \mathbb{N}_0$ with $c > 0$ such that $n = 2^c 3^d$ and $x = d + c\eta$;
- (v) $x \in \mathbb{Q} + \mathbb{Q}\eta$.

If these conditions fail, then $\eta \log 2 \notin \mathcal{L}$ and E_2 is log-strongly transcendental. If they hold, then

$$E_2^c = \frac{m}{2^d} \in \mathbb{Q}_{>0}. \quad (39)$$

Proof. Interchange the roles of 2 and 3, of m and n , and of ϑ and η in the proof of Theorem 17.2. For the strong step, use the $\mathbb{Q}$ -linear independence of $\log 3, \log n, \log 2$ when n is not three-smooth. $\square$

Corollary 17.4 (The two branches are mutually exclusive). *For a nonintegral solution, E_3 and E_2 cannot both be algebraic. Equivalently, at least one of $\vartheta \log 3$ and $\eta \log 2$ lies outside $\mathcal{L}$ under a counterexample as well as unconditionally.*

Proof. By Theorems 17.2 and 17.3, algebraicity of E_3 is equivalent to three-smoothness of m and algebraicity of E_2 to three-smoothness of n . Proposition 16.4 forbids both, since eliminating x from $x = a + b\vartheta$ and $x = d + c\eta$ produces the quadratic (32) for the transcendental ϑ . $\square$

So under a counterexample the unconditional dichotomy of Theorem 17.1 becomes a genuine classification rather than an alternative: at most one of the two constants is algebraic, and each is algebraic exactly in its own smooth-output branch, with an explicit radical description.

Smooth branch	Return kernel	Symmetric constants
$m = 2^a 3^b, b > 0$	$\{2^r : r \in \mathbb{Q}\}$	$E_3^b = n/3^a \in \mathbb{Q}_{>0}$; E_2 log-strongly transcendental
$n = 2^c 3^d, c > 0$	$\{3^s : s \in \mathbb{Q}\}$	$E_2^c = m/2^d \in \mathbb{Q}_{>0}$; E_3 log-strongly transcendental
neither output smooth	no coordinate axis	both E_3 and E_2 log-strongly transcendental

The three rows of the table are exactly the three surviving rows of the primitive-core trichotomy of Section 9: the first is the case $w = 3$, in which the canonical radical $E = w^\vartheta$ of Section 6 degenerates to E_3 itself with index $d = b$; the second is its mirror $v = 2$; and the third is the generic case with two independent external prime supports. The missing fourth pattern, both outputs three-smooth, is the one excluded by Proposition 16.4.

Two remarks are worth making about the direction of the implications. First, nothing here proves that E_3 or E_2 is transcendental: a counterexample would make one of them an explicit algebraic radical, and Corollary 17.4 only says that it cannot make both algebraic. The conjecture remains open, and the smallest instance — irrationality of $E_3 = 3^{\log_2 3}$ — is exactly the wall named in Section 6. Second, the implication that is usable in the other direction is the strong one: proving $\vartheta \log 3 \notin \mathcal{L}$ would kill the entire $w = 3$ branch, and proving $\eta \log 2 \notin \mathcal{L}$ would kill the $v = 2$ branch, leaving only the generic case, in which both outputs carry prime support outside $\{2, 3\}$ and the local determinant methods have two independent external primes to work with.

Status: [\[paper proof in this report\]](#) — paper proofs in this report. Theorems 17.2 and 17.3 use the six exponentials theorem for the algebraicity equivalences and Baker’s theorem together with Roy’s strong six exponentials theorem for the $\mathcal{L}$ -statements; none of those classical inputs is formalized in the repository. Corollary 17.4 adds only Proposition 16.4, whose sole nonelementary ingredient is the kernel-verified transcendence of ϑ , so the mutual exclusion of the two smooth branches — as opposed to the transcendence conclusions attached to them — is within immediate reach of the Lean development.

18 Finite-difference obstructions

The first local obstruction recorded in the corpus forbids three candidates in arithmetic progression under the convenient hypothesis $h^5 \leq n$. That hypothesis is an artifact of the estimate, not of the mechanism: the sharp arbitrary-order curvature theorem is now kernel-verified for every $r \geq 2$, and it contains the three- and four-term cases as specializations.

For a function f and a step $h > 0$ write $\Delta_h f(t) = f(t+h) - f(t)$ and $\Delta_h^r = \Delta_h \circ \cdots \circ \Delta_h$ (r factors). Iterating the fundamental theorem of calculus gives the mean-value identity in integral form,

$$\Delta_h^r f(n) = \int_{[0,h]^r} f^{(r)}(n + t_1 + \cdots + t_r) dt_1 \cdots dt_r. \quad (40)$$

For $f(t) = t^\vartheta$ and $r \geq 2$ the derivative is

$$f^{(r)}(t) = (\vartheta)_{\underline{r}} t^{\vartheta-r}, \quad (\vartheta)_{\underline{r}} := \prod_{j=0}^{r-1} (\vartheta - j) = \vartheta(\vartheta-1) \cdots (\vartheta-r+1).$$

Since $1 < \vartheta < 2$ the falling factorial never vanishes, so $f^{(r)}$ has a fixed nonzero sign on $(0, \infty)$, and $|f^{(r)}|$ decreases in t . These two facts — nonvanishing and decay — are the whole content of the obstruction.

Theorem 18.1 (Higher-difference progression obstruction). *Let $r \geq 2$, let $n, h \in \mathbb{N}$ be positive, and set $c_r = |(\vartheta)_{\underline{r}}|$. If all of $n, n+h, \dots, n+rh$ are candidates, then*

$$1 \leq c_r h^r n^{\vartheta-r}.$$

Such a progression is therefore impossible whenever

$$c_r h^r < n^{r-\vartheta}. \quad (41)$$

Proof. If all $n + jh$ are candidates, every $(n + jh)^\vartheta$ is an integer, so the integer combination $\Delta_h^r f(n) = \sum_{j=0}^r (-1)^{r-j} \binom{r}{j} f(n + jh)$ is an integer. By (40) and the fixed sign of $f^{(r)}$ it is nonzero, hence of absolute value at least one. Bounding $|f^{(r)}|$ by its value at the left endpoint n ,

$$1 \leq |\Delta_h^r f(n)| \leq c_r h^r n^{\vartheta-r}. \quad \square$$

Status: [\[kernel-verified Lean\]](#) — `HigherDifference` proves the generic iterated mean-value identity `exists_iterated_fwdDiff_eq_pow_mul`, its real-power specialization `exists_iterated_fwdDiff_rpow_point`, integrality of the binomial forward difference, and the obstruction `not_arithmetic_progression_twoBaseNaturalCandidates_of_curvature` uniformly for every $r \geq 2$. The displayed iterated-integral identity (40) is not used by the Lean proof; the formal argument iterates the mean-value theorem instead, which avoids any integration theory.

18.1 The sharp three-term criterion

For $r = 2$ the criterion (41) reads

$$\vartheta(\vartheta - 1)h^2 < n^{2-\vartheta}, \quad \text{i.e.} \quad h < 1.038547802 \dots \cdot n^{0.2075187\dots}, \quad (42)$$

strictly weaker as a hypothesis than $h \leq n^{1/5}$, whose exponent is only 0.2. The real form is `second_difference_logThreeDivLogTwo_ap_of_curvature` in `SharperProgressions`, built on the curvature estimate `second_difference_logThreeDivLogTwo_ap_bound`.

Almost all of the gain survives in a purely rational hypothesis, which is what a kernel can check without any real-number decision procedure. The step exponent $83/400 = 0.2075$ is admissible:

Theorem 18.2 (Sharp rational three-term criterion). *Let $n, h \geq 1$ be integers with $h^{400} \leq n^{83}$. Then $n, n + h, n + 2h$ are not all candidates.*

Status: [\[kernel-verified Lean\]](#) — `SharpProgression` proves `second_difference_logThreeDivLogTwo_ap_sharp` and `not_three_arithmetic_progression_twoBaseNaturalCandidates_sharp`. The two numerical ingredients are the enclosure $\vartheta < 317/200$, replayed by the kernel from the exact power inequality $3^{200} < 2^{317}$ through `logThreeDivLogTwo_lt_317_div_200`, and the consequence $\vartheta(\vartheta - 1) < 37089/40000 < 1$. For $r = 3$, `ThirdDifference` kernel-checks the exact third-difference identity and the corresponding four-candidate obstruction.

18.2 Thresholds for higher order

Writing (41) as a bound on the step, a progression of length $r + 1$ starting at n is impossible once

$$h < C_r n^{1-\vartheta/r}, \quad C_r = c_r^{-1/r}.$$

Table 2 lists the first thresholds.

18.3 Why longer progressions do not close the problem

As r grows the admissible step exponent $1 - \vartheta/r$ approaches one, which looks tantalizing: for large r the theorem forbids progressions with steps almost as large as the starting point itself. It does not contradict the exact candidate monoid, and it cannot be bootstrapped into a proof, for a purely combinatorial reason.

r	length $r + 1$	exponent $1 - \vartheta/r$	C_r
2	3	0.207518750	1.038547802
3	4	0.471679166	1.374849673
4	5	0.603759375	1.164124485
5	6	0.683007500	0.946705202
6	7	0.735839583	0.778537835
7	8	0.773576786	0.652646038
8	9	0.801879687	0.557368999

Table 2: Step thresholds for the higher-difference obstruction. A progression of length $r + 1$ with step $h < C_r n^{1-\vartheta/r}$ cannot consist of candidates.

Observation 18.3 (Combinatorial limitation). Conditionally on the structure described in Section 5, a dyadic block near X contains $O(\log X)$ candidates, so the average additive gap inside the block is of order $X/\log X$. For each fixed r this exceeds $X^{1-\vartheta/r}$ once X is large. A rank-two multiplicative monoid carries no reason whatsoever to contain long *additive* progressions; the hypothesis of Theorem 18.1 is therefore vacuous for the configurations that actually occur.

Extracting leverage from long progressions would require r growing with the height, control of the rapidly growing constant $c_r = |(\vartheta)_r|$, and — decisively — a combinatorial source guaranteeing such progressions exist. None of the three is available. The productive response is to drop the additive-structure hypothesis altogether, which is what Section 19 does.

18.4 The optimized variable-order scale

Theorem 18.1 is uniform in r , so nothing forces the difference order to stay fixed while the starting point grows. Table 2 records the first fixed-order thresholds, but each row is only a snapshot of a one-parameter family that may be re-optimized at every height. Carrying out that optimization strengthens every fixed-order corollary at once and, more valuably, replaces the vague complaint that the candidate density “feels too small” by a single explicit constant. Throughout this subsection $\log$ denotes the natural logarithm.

The exact derivative coefficient. Fix $r \geq 2$ and recall the constant $c_r = |(\vartheta)_r|$ of Theorem 18.1. Since $1 < \vartheta < 2$, the first two factors ϑ and $\vartheta - 1$ are positive while every later factor $\vartheta - j$ with $j \geq 2$ is negative, so taking absolute values termwise gives

$$c_r = \vartheta(\vartheta - 1) \prod_{j=2}^{r-1} (j - \vartheta),$$

and the recursion $\Gamma(z + 1) = z\Gamma(z)$ telescopes the product into closed form:

$$c_r = \frac{\vartheta(\vartheta - 1)}{\Gamma(2 - \vartheta)} \Gamma(r - \vartheta). \quad (43)$$

So c_r grows at factorial speed, which is exactly why no fixed order can be optimal for all n : the constant eventually defeats the improving exponent. Stirling’s formula applied to (43) gives the root asymptotic that the optimization needs,

$$c_r^{1/r} = \frac{r}{e} \left(1 + O\left(\frac{\log r}{r}\right) \right). \quad (44)$$

Optimizing the step threshold over the order. The formalized condition (41) forbids a progression of length $r + 1$ as soon as $h < c_r^{-1/r} n^{1-\vartheta/r}$, and substituting (44) rewrites that threshold as

$$c_r^{-1/r} n^{1-\vartheta/r} = \frac{en}{r} \exp\left(-\frac{\vartheta \log n}{r}\right) \left(1 + O\left(\frac{\log r}{r}\right)\right). \quad (45)$$

The two r -dependent factors pull in opposite directions: $c_r^{-1/r} \sim e/r$ decays in r , while $n^{-\vartheta/r}$ increases to 1. Discarding the lower-order factor and differentiating the logarithm of the leading term,

$$\frac{d}{dr} \left(1 + \log n - \log r - \frac{\vartheta \log n}{r}\right) = -\frac{1}{r} + \frac{\vartheta \log n}{r^2},$$

which is positive for $r < \vartheta \log n$, vanishes at $r = \vartheta \log n$, and is negative afterwards. The leading term is therefore unimodal in real $r > 0$ with a unique maximum at $r = \vartheta \log n$, and at that point the explicit factor e cancels against $\exp(-\vartheta \log n/r) = e^{-1}$, leaving the clean scale $n/(\vartheta \log n)$.

Theorem 18.4 (Optimized variable-order obstruction). *Let $0 < \varepsilon < 1$. For all sufficiently large n choose an integer*

$$r_n = \vartheta \log n + O(1), \quad r_n \geq 2.$$

If

$$1 \leq h \leq (1 - \varepsilon) \frac{n}{\vartheta \log n},$$

then $n, n + h, \dots, n + r_n h$ are not all candidates. Moreover, the largest leading-order step scale obtainable by optimizing the curvature condition (41) over r is

$$h_{\max}(n) \sim \frac{n}{\vartheta \log n} = \frac{\log 2}{\log 3} \cdot \frac{n}{\log n} = 0.6309297535 \dots \cdot \frac{n}{\log n}.$$

Proof. Take r_n to be either nearest integer to $\vartheta \log n$. Then $r_n \rightarrow \infty$, so (44) and hence (45) apply, and the quotient of $c_{r_n}^{-1/r_n} n^{1-\vartheta/r_n}$ by $n/(\vartheta \log n)$ tends to 1. The fixed margin $1 - \varepsilon$ absorbs both the Stirling error $1 + O(\log r/r)$ and the $O(1)$ rounding of $\vartheta \log n$ to an integer, so (41) holds for every h in the stated range, and Theorem 18.1 excludes the progression.

For the optimality assertion, the leading approximation in (45) is unimodal in real $r > 0$ with its unique maximum at $r = \vartheta \log n$, where its value is $n/(\vartheta \log n)$. Any regime in which r stays bounded, or grows outside a constant multiple of $\log n$, returns a strictly smaller order or a strictly smaller leading constant. $\square$

Status: [\[paper proof in this report\]](#) — the obstruction being optimized is [\[kernel-verified Lean\]](#) and holds uniformly for every $r \geq 2$ with no upper bound on the order, so substituting $r = r_n$ inside the formalized statement is legitimate; the relevant module is `HigherDifference`. The optimization itself, the Gamma evaluation (43), and the Stirling asymptotic (44) are paper arguments and have not been submitted to the kernel.

Comparison with the conditional candidate spacing. A dyadic block $[X, 2X)$ is an interval of length one in exponent space, so the exact block count of Section 5 applies verbatim: conditional on a counterexample with least nonintegral generator β ,

$$\#(\mathcal{C} \cap [X, 2X)) = \frac{\log_2 X}{\beta} + O(1) = \frac{\log X}{\beta \log 2} + O_\beta(1). \quad (46)$$

The average additive spacing between consecutive candidates inside that block is therefore

$$\sim \beta \log 2 \cdot \frac{X}{\log X},$$

whereas Theorem 18.4 taken at $n \asymp X$ reaches only

$$h_{\max} \sim \frac{1}{\vartheta} \cdot \frac{X}{\log X}.$$

Both are of order $X/\log X$, so the comparison is a pure constant, and in it the transcendental ϑ cancels against the $\log 2$:

$$\frac{\beta \log 2}{1/\vartheta} = \beta \vartheta \log 2 = \beta \log 3. \quad (47)$$

The kernel-verified zero-free region $2^x < 4096$ forces $\beta > 12$ (Section 5), which already makes the ratio exceed $12 \log 3 = 13.1833\dots$; the interval scan of Section 25 raises it past $20 \log 3 = 21.9722\dots$, and the extended directed-rounding scan through 2^{27} raises it past $27 \log 3 = 29.6625\dots$. The last two improvements are [software computation] and do not enter any kernel-verified statement; the exponent bound the kernel actually certifies remains $x \geq 12$.

The same constant governs a second, purely cardinal comparison, which is the more telling of the two. The configuration forbidden by Theorem 18.4 consists of

$$r_n + 1 \sim \vartheta \log X$$

candidates arranged in one structured pattern, while the whole dyadic block contains only $\sim \log X/(\beta \log 2)$ candidates by (46). The ratio is again exactly $\beta \log 3$: the optimized theorem asks a block to supply about $\beta \log 3$ times as many candidates as it can possibly hold. Concretely, at $X = 2^{100}$ the block $[X, 2X)$ holds $100/\beta + O(1)$ candidates — about eight if β were as small as its kernel-certified floor, and about three under the 2^{27} scan — while the shortest progression the theorem forbids at that height already needs $r_n + 1 \approx 111$ terms.

Remark 18.5 (Quantified barrier). Optimizing over r materially strengthens the local theorem: the admissible step scale rises from $n^{1-\vartheta/r}$ at fixed r to $n^{1-o(1)}$, missing n itself only by the single logarithmic factor $\vartheta \log n$. It still cannot contradict the conditional monoid, because the conditional candidate density in a dyadic block is too small by the fixed factor $\beta \log 3 > 13.18$. No argument using only the number of candidates in a dyadic block can force the progression that the current finite-difference functional requires. A successful continuation must exploit additional arithmetic correlation — for example valuation congruences that force a progression of a special shape — or replace Δ_h^r by a different integer-valued functional with a better curvature-to-spacing constant.

This is a limitation of the method, not evidence against Conjecture 2.1. Its value is to replace “the density seems too weak” by a sharp asymptotic constant and a concrete numerical target that any improved functional must beat.

Agreement with the one-logarithm deficit. Section 22 audits the semigroup determinant in the same currency and identifies the same bottleneck. In both analyses the binding quantity is the conditional dyadic population $\asymp \log X$ of (46), and in both the archimedean inequality turns out to be comfortably satisfied by the exact monoid rather than violated by it. To that extent the two diagnoses agree. They differ, instructively, on the size of the shortfall. The semigroup determinant needs $r \asymp (\log X)^2$ nodes and is supplied $\asymp \log X$: a deficit of a full power of $\log X$. The optimized

finite difference needs $\vartheta \log X$ nodes and is supplied $\log X/(\beta \log 2)$: a deficit of the bounded factor $\beta \log 3$. The finite difference is thus quantitatively much closer — but only because it buys that closeness with a hypothesis it cannot discharge, namely that the candidates lie in exact arithmetic progression, which Observation 18.3 says a rank-two multiplicative monoid has no reason to provide. Discarding that hypothesis is exactly what the determinants of Sections 19 and 20 do, and the one-logarithm deficit is the precise price charged for it. Read together, the two sections say that the present archimedean toolkit is short by one factor of $\log X$ in generality or by a factor of $\beta \log 3$ in strength, and that tuning constants converts neither shortfall into the other.

19 New result I: arbitrary-node lattice repulsion

The higher-difference theorem of Section 18 uses equally spaced nodes, because the forward-difference operator is built from an arithmetic progression. The same integer-versus-small-real principle extends to *arbitrary* nodes through interpolation determinants. The resulting statement constrains every configuration of candidates, not only the additively structured ones, which is exactly the gap identified in Observation 18.3.

19.1 The determinant identity

For distinct real numbers $x_0 < \dots < x_r$ and a function f , write $f[x_0, \dots, x_r]$ for the divided difference. The standard determinant identity is

$$\det \begin{pmatrix} 1 & x_0 & \dots & x_0^{r-1} & f(x_0) \\ 1 & x_1 & \dots & x_1^{r-1} & f(x_1) \\ \vdots & \vdots & & \vdots & \vdots \\ 1 & x_r & \dots & x_r^{r-1} & f(x_r) \end{pmatrix} = V(x_0, \dots, x_r) f[x_0, \dots, x_r], \quad (48)$$

where

$$V(x_0, \dots, x_r) = \prod_{0 \leq i < j \leq r} (x_j - x_i).$$

If $f \in C^r$ on the interval spanned by the nodes, the generalized mean-value theorem for divided differences supplies some $\xi \in (x_0, x_r)$ with

$$f[x_0, \dots, x_r] = \frac{f^{(r)}(\xi)}{r!}. \quad (49)$$

For $f(t) = t^\vartheta$ this is again $f^{(r)}(t) = (\vartheta)_{\underline{r}} t^{\vartheta-r}$, and since ϑ is irrational — in particular not an integer — the coefficient $(\vartheta)_{\underline{r}}$ never vanishes.

19.2 The arbitrary-node theorem

Theorem 19.1 (Vandermonde repulsion for candidates). *Let $r \geq 2$ and let*

$$1 \leq m_0 < m_1 < \dots < m_r$$

be natural candidates, so that $m_i^\vartheta \in \mathbb{N}$ for every i . Then

$$\prod_{0 \leq i < j \leq r} (m_j - m_i) \geq \frac{r!}{|(\vartheta)_{\underline{r}}|} m_0^{r-\vartheta}. \quad (50)$$

Proof. Apply (48) to $f(t) = t^\vartheta$ with the nodes m_i . Every entry of the matrix is an integer: the ordinary powers m_i^k are integral by construction, and the last column is integral precisely because the m_i are candidates. By (49) the determinant equals

$$D = V(m_0, \dots, m_r) \frac{(\vartheta)_r}{r!} \xi^{\vartheta-r} \quad (m_0 < \xi < m_r).$$

Each factor is nonzero, so $D \in \mathbb{Z} \setminus \{0\}$ and therefore $|D| \geq 1$. Because $r \geq 2 > \vartheta$ the exponent $\vartheta - r$ is negative, so $\xi^{\vartheta-r} \leq m_0^{\vartheta-r}$. Therefore

$$1 \leq |D| \leq V(m_0, \dots, m_r) \frac{|(\vartheta)_r|}{r!} m_0^{\vartheta-r},$$

which rearranges to (50). □

Status: [\[paper proof in this report\]](#) for the unconditional paper theorem;

[\[kernel-verified Lean\]](#) for the formal determinant and repulsion chain.

`ArbitraryNodeDeterminant` proves integrality, the Vandermonde identities, and the hypothesis-carrying repulsion step, while `ArbitraryNodeRepulsion` packages the final bound under the explicit hypothesis `RpowDividedDifferenceMeanValue`. The corpus does not assert that analytic hypothesis.

Remark 19.2 (Formalization state). The accepted module `ArbitraryNodeDeterminant` supplies the algebraic half deliberately free of analysis — the node matrix, an explicit integral model, integrality of the determinant with the resulting bound $1 \leq |\det|$, the Vandermonde product and its positivity on strictly increasing nodes. The later endpoint theorem is also accepted, but its divided-difference mean-value statement for $t \mapsto t^\vartheta$ remains a visible hypothesis rather than a proved ingredient. Thus the conditional Lean implication may be cited as machine-checked; the unconditional analytic discharge is supplied only by the paper proof here.

19.3 Three points: a lattice-triangle inequality

For $r = 2$ the determinant is

$$D = \det \begin{pmatrix} 1 & m_0 & m_0^\vartheta \\ 1 & m_1 & m_1^\vartheta \\ 1 & m_2 & m_2^\vartheta \end{pmatrix},$$

and $|D|$ is twice the area of the lattice triangle with vertices (m_i, m_i^ϑ) . Strict convexity of $t \mapsto t^\vartheta$ makes the area positive; lattice integrality makes it at least $1/2$. The bound below is the quantitative form of that picture.

Corollary 19.3 (Three-candidate span). *If three candidates lie in $[N, N + H]$ with $N \geq 1$, then*

$$H \geq \left(\frac{8}{\vartheta(\vartheta-1)} \right)^{1/3} N^{(2-\vartheta)/3}. \quad (51)$$

Rounded to ten decimal places the constant is 2.0510723957 and the exponent 0.1383458331; in particular $H \geq 2.0510723956 N^{0.1383458330}$.

Proof. Write $a = m_1 - m_0$ and $b = m_2 - m_1$. Then $a, b > 0$ and $a + b \leq H$, so the Vandermonde product satisfies

$$V = ab(a + b) \leq \frac{H^3}{4},$$

the continuous maximum being attained at $a = b = H/2$. Theorem 19.1 at $r = 2$ gives

$$V \geq \frac{2}{\vartheta(\vartheta - 1)} N^{2-\vartheta}.$$

Combining the two inequalities yields (51). □

19.4 General local packing bound

Put

$$R_r := \binom{r+1}{2} = \frac{r(r+1)}{2}, \quad C_r := \left(\frac{r!}{|(\vartheta)_r|} \right)^{1/R_r}, \quad \alpha_r := \frac{r - \vartheta}{R_r} = \frac{2(r - \vartheta)}{r(r+1)}.$$

Every pairwise gap among points of $[N, N + H]$ is at most H , so the Vandermonde product is at most H^{R_r} , and Theorem 19.1 immediately gives the following.

Corollary 19.4 (At most r points below the repulsion scale). *If $r + 1$ candidates lie in $[N, N + H]$, then*

$$H \geq C_r N^{\alpha_r}. \tag{52}$$

Equivalently, every interval of length strictly less than $C_r N^{\alpha_r}$ beginning at or after N contains at most r candidates.

The exponent can be optimized in r . As a function of a real variable $r > \vartheta$,

$$\alpha(r) = \frac{2(r - \vartheta)}{r(r + 1)}$$

attains its maximum at

$$r_* = \vartheta + \sqrt{\vartheta^2 + \vartheta} = 3.6090841941\dots,$$

so the best integer order in this ordinary-power family is $r = 4$ — not the three-point case that geometry makes most vivid. Table 3 lists the constants.

Corollary 19.5 (Best ordinary-power local statement). *Every interval*

$$[N, N + H], \quad H < 1.4602287460 N^{0.2415037499},$$

contains at most four natural candidates. The exact constant is $C_4 = (4! / |(\vartheta)_4|)^{1/10} = 1.4602287461$ to ten decimal places, and the exact exponent is $\alpha_4 = (4 - \vartheta)/10 = 0.2415037499\dots$

r	R_r	$ (\vartheta)_r $	α_r	C_r
2	3	0.9271436280	0.1383458331	1.2920946431
3	6	0.3847993728	0.2358395832	1.5805909191
4	10	0.5445055422	0.2415037499	1.4602287461
5	15	1.3150013031	0.2276691666	1.3510881062
6	21	4.4907787617	0.2102398809	1.2735045942
7	28	19.8269566337	0.1933941964	1.2187063909
8	36	107.3637136680	0.1781954861	1.1790125192

Table 3: Ordinary-power arbitrary-node repulsion constants. For $r = 2$ the sharper three-point constant of Corollary 19.3 uses the exact maximum of the three-gap product rather than the crude bound H^{R_2} .

19.5 A bound for any longer interval

The span estimate converts into a counting inequality. Let $m_1 < \dots < m_s$ be all candidates in $[N, N + H]$. For each $i \leq s - r$ the consecutive window $m_i, \dots, m_{i+r}$ has span at least $C_r N^{\alpha_r}$ by Corollary 19.4. Summing these spans, each adjacent gap $m_{j+1} - m_j$ is counted at most r times, so

$$(s - r)C_r N^{\alpha_r} \leq r(m_s - m_1) \leq rH.$$

Corollary 19.6 (Arbitrary-node counting inequality). *For every $r \geq 2$,*

$$\#(\mathcal{C} \cap [N, N + H]) \leq r + \frac{r}{C_r} H N^{-\alpha_r}. \quad (53)$$

This is a genuine strengthening of progression avoidance: it constrains every possible configuration of candidates in an interval, with no additive hypothesis at all. It remains far from the conditional benchmark of Section 5. At the optimal $r = 4$ and $H = N$, the right-hand side of (53) is $O(N^{0.7585\dots})$, whereas the exact population of a failed conjecture would be only $O(\log N)$. The gap is enormous, and closing it by ordinary powers alone is hopeless — which motivates enlarging the column space, the subject of Section 20.

19.6 Relation to higher differences

For equally spaced nodes $m_i = n + ih$ the Vandermonde factor is explicit:

$$V = h^{R_r} \prod_{j=1}^r j!.$$

Substituting in (48) recovers an integer interpolation remainder closely related to the r th forward difference of Section 18. The committed repository theorem is sharper for this special configuration, because it uses the exact binomial combination rather than bounding all pairwise gaps by a common span. Conversely, Theorem 19.1 applies when the nodes carry no additive structure whatsoever. The two results are complementary, and only the arbitrary-node form escapes the vacuity diagnosed in Observation 18.3.

20 New result II: semigroup generalized Vandermonde determinants

The determinant of Section 19 used only the integral columns $1, m, m^2, \dots$ and one extra integral column m^ϑ . But a candidate carries much more integral data. If $m^\vartheta = n \in \mathbb{N}$, then for every $a, b \in \mathbb{N}_0$,

$$m^{a+b\vartheta} = m^a n^b \in \mathbb{N}. \quad (54)$$

The full set of available exponents is therefore

$$\Lambda_\vartheta := \{a + b\vartheta : a, b \in \mathbb{N}_0\}.$$

Since ϑ is irrational, all pairs (a, b) give distinct exponents.

20.1 Positivity of generalized Vandermonde determinants

Let

$$0 \leq \lambda_0 < \lambda_1 < \dots < \lambda_r \quad \text{and} \quad 0 < x_0 < x_1 < \dots < x_r.$$

The functions x^{λ_j} form an extended complete Chebyshev system on $(0, \infty)$. Their Wronskian is classical in the Karlin–Studden theory of Tchebycheff systems:

$$W(x) = x^{\sum_j \lambda_j - R_r} \prod_{0 \leq i < j \leq r} (\lambda_j - \lambda_i) > 0. \quad (55)$$

Indeed, after factoring x^{λ_j} from column j and x^{-i} from derivative row i , the remaining determinant is the Vandermonde determinant of the λ_j , expressed in the falling-factorial basis. Positivity of all initial Wronskians implies strict positivity of every ordered evaluation determinant.

For completeness, there is also an elementary zero-count proof. Put $x = e^t$. A nonzero linear combination of x^{λ_j} becomes an exponential polynomial $\sum c_j e^{\lambda_j t}$. After factoring the lowest exponential, differentiation removes the constant term and leaves one fewer exponential. Induction with Rolle’s theorem shows that a combination of $r + 1$ such functions has at most r real zeros counting multiplicity. Consequently an evaluation determinant at $r + 1$ ordered nodes cannot vanish. Its sign is constant on the connected chamber $x_0 < \dots < x_r$ and is positive in the coalescing-node limit by (55).

Lemma 20.1 (Generalized Vandermonde positivity). *For strictly increasing nonnegative exponents λ_j and strictly increasing positive nodes x_i ,*

$$\det(x_i^{\lambda_j})_{0 \leq i, j \leq r} > 0. \quad (56)$$

The two arguments just sketched are carried out in full, with all induction bookkeeping made explicit, in Appendix D.

20.2 Integer determinants from candidates

Theorem 20.2 (Semigroup integer determinant). *Let $m_0 < \dots < m_r$ be natural candidates. Choose distinct pairs $(a_j, b_j) \in \mathbb{N}_0^2$ and order*

$$\lambda_j = a_j + b_j \vartheta$$

so that $\lambda_0 < \dots < \lambda_r$. Then

$$D_{\lambda}(\mathbf{m}) := \det(m_i^{\lambda_j})_{0 \leq i, j \leq r} \quad (57)$$

is a positive integer. In particular, $D_{\lambda}(\mathbf{m}) \geq 1$.

Proof. By (54), each entry equals

$$m_i^{a_j} (m_i^{\vartheta})^{b_j} \in \mathbb{N},$$

so the determinant is an integer. It is strictly positive by Lemma 20.1. $\square$

Status: [\[paper proof in this report\]](#)

This theorem packages all mixed integrality consequences into one reusable object. It is not the same as the repository’s exponential interpolation determinant in `MultiplicativeRank.lean`: here the rows are arbitrary natural candidates and the columns are chosen from the two-dimensional exponent semigroup.

20.3 Divided-difference factorization

Let $f_j(x) = x^{\lambda_j}$ and let

$$V(\mathbf{m}) = \prod_{i < j} (m_j - m_i).$$

Applying the Newton divided-difference row transformation gives

$$D_{\lambda}(\mathbf{m}) = V(\mathbf{m}) \det(f_j[m_0, \dots, m_i])_{0 \leq i, j \leq r}. \quad (58)$$

The transformation is exact; its determinant is $V(\mathbf{m})^{-1}$.

Suppose

$$N \leq m_0 < \dots < m_r \leq N + H \leq 2N.$$

For each divided difference, the mean-value theorem gives a point in $[N, 2N]$ and

$$|f_j[m_0, \dots, m_i]| \leq \frac{|(\lambda_j)_i|}{i!} \sup_{N \leq x \leq 2N} x^{\lambda_j - i}. \quad (59)$$

We need a uniform elementary coefficient estimate.

Lemma 20.3 (Falling-factorial coefficient bound). *For $\lambda \geq 0$ and $i \in \mathbb{N}_0$,*

$$\frac{|(\lambda)_i|}{i!} \leq 2^{\lambda+i}. \quad (60)$$

Proof. If $0 \leq \lambda \leq 1$, each factor satisfies $|\lambda - k| \leq k + 1$, so the quotient is at most 1. If $\lambda \geq 1$, let $n = \lceil \lambda \rceil$. Then

$$\frac{|(\lambda)_i|}{i!} \leq \prod_{k=0}^{i-1} \frac{\lambda + k}{k + 1} \leq \prod_{k=0}^{i-1} \frac{n + k}{k + 1} = \binom{n + i - 1}{i}.$$

The last binomial coefficient is at most the sum of its row, 2^{n+i-1} , and $n - 1 \leq \lambda$. Hence it is at most $2^{\lambda+i}$. $\square$

For $x \in [N, 2N]$,

$$x^{\lambda-i} \leq 2^\lambda N^{\lambda-i}. \quad (61)$$

If $\lambda - i \geq 0$, use $x \leq 2N$; if it is negative, use $x \geq N$. Combining (59), (60), and (61),

$$|f_j[m_0, \dots, m_i]| \leq 2^{2\lambda_j+i} N^{\lambda_j-i}. \quad (62)$$

20.4 Quantitative semigroup gap theorem

Set

$$R := R_r = \frac{r(r+1)}{2}, \quad \Sigma := \sum_{j=0}^r \lambda_j.$$

By the Leibniz formula and (62), every permutation term in the small determinant in (58) is at most

$$2^{2\Sigma+R} N^{\Sigma-R},$$

and there are $(r+1)!$ terms. Since every pairwise node difference is at most H , $V(\mathbf{m}) \leq H^R$. The positive integer lower bound from Theorem 20.2 now gives the following.

Theorem 20.4 (Semigroup determinant gap). *Let $r \geq 1$. Let $\lambda_0 < \dots < \lambda_r$ be any $r+1$ exponents in Λ_ϑ , and put $R = r(r+1)/2$ and $\Sigma = \sum_j \lambda_j$. If $r+1$ natural candidates lie in*

$$[N, N+H] \subseteq [N, 2N], \quad N \geq 1,$$

then

$$H \geq ((r+1)!)^{-1/R} 2^{-1-2\Sigma/R} N^{1-\Sigma/R}. \quad (63)$$

Proof. Equations (58) and (62) imply

$$1 \leq D_\lambda(\mathbf{m}) \leq H^R (r+1)! 2^{R+2\Sigma} N^{\Sigma-R}.$$

Taking R th roots and rearranging proves (63). $\square$

The exponent $1 - \Sigma/R$ is the key. We should choose the $r+1$ smallest available semigroup exponents to minimize Σ .

20.5 The ordered exponent semigroup

Write the elements of Λ_ϑ in increasing order:

$$0 = \lambda_0 < \lambda_1 < \lambda_2 < \dots.$$

The beginning is

$$0, 1, \vartheta, 2, 1+\vartheta, 3, 2\vartheta, 2+\vartheta, 4, 1+2\vartheta, 3+\vartheta, 3\vartheta, 5, \dots$$

The counting function is elementary:

$$\begin{aligned} L(T) &:= \#\{\lambda \in \Lambda_\vartheta : \lambda \leq T\} \\ &= \sum_{0 \leq b \leq T/\vartheta} (\lfloor T - b\vartheta \rfloor + 1) = \frac{T^2}{2\vartheta} + O(T). \end{aligned} \quad (64)$$

Irrationality of ϑ ensures that different pairs (a, b) are counted distinctly. Inverting (64) and summing the order statistics gives

$$\lambda_r = \sqrt{2\vartheta r} + O(1), \quad (65)$$

$$\Sigma_r := \sum_{j=0}^r \lambda_j = \frac{2}{3} \sqrt{2\vartheta} r^{3/2} + O(r), \quad (66)$$

$$\frac{\Sigma_r}{R_r} = \frac{4}{3} \sqrt{\frac{2\vartheta}{r}} + O(r^{-1}). \quad (67)$$

Consequently the power of N in (63) tends to 1:

$$1 - \frac{\Sigma_r}{R_r} = 1 - \frac{4}{3} \sqrt{\frac{2\vartheta}{r}} + O(r^{-1}). \quad (68)$$

r	Σ_r	$1 - \Sigma_r/R_r$
2	2.5849625007	0.1383458331
3	4.5849625007	0.2358395832
4	7.1699250014	0.2830074999
5	10.1699250014	0.3220049999
6	13.3398500029	0.3647690475
7	16.9248125036	0.3955424106
8	20.9248125036	0.4187552082
9	25.0947375050	0.4423391666
10	29.6797000058	0.4603690908
12	39.4345875079	0.4944283653
20	87.7939875195	0.5819333928

Table 4: Exponent gain from using the first $r + 1$ mixed exponents in $\mathbb{N}_0 + \mathbb{N}_0\vartheta$. The ordinary-power family peaks and then decays; the semigroup family improves monotonically in these data and tends to exponent 1.

20.6 An explicit coarse form

For the global counting corollary we need no delicate lattice-point error term. A simple bound suffices.

Lemma 20.5 (Crude semigroup weight bound). *For $r \geq 2$, the first $r + 1$ semigroup exponents satisfy*

$$\lambda_r \leq 6\sqrt{r}, \quad \frac{\Sigma_r}{R_r} \leq \frac{12}{\sqrt{r}}. \quad (69)$$

Proof. Let $q = \lceil \sqrt{r+1} \rceil$. The $(q+1)^2$ pairs $0 \leq a, b \leq q$ give distinct exponents and each is at most $q(1+\vartheta) < 3q$. Since $(q+1)^2 > r+1$, the r th ordered exponent is below $3q$. For $r \geq 2$, $q \leq \sqrt{r+1} + 1 \leq 2\sqrt{r}$, hence $\lambda_r \leq 6\sqrt{r}$. Then

$$\Sigma_r \leq (r+1)\lambda_r \leq 6(r+1)\sqrt{r},$$

and division by $R_r = r(r+1)/2$ gives the second inequality. $\square$

For $r \geq 144$, (69) gives $\Sigma_r/R_r \leq 1$. Also

$$((r+1)!)^{1/R_r} \leq (r+1)^{(r+1)/R_r} = (r+1)^{2/r} \leq 3$$

for $r \geq 2$. Therefore Theorem 20.4 yields a particularly clean form.

Corollary 20.6 (Explicit near-linear span). *If $r \geq 144$ and $r+1$ natural candidates lie in $[N, N+H] \subseteq [N, 2N]$, then*

$$H \geq \frac{1}{24} N^{1-12/\sqrt{r}}. \quad (70)$$

Proof. In (63), use $((r+1)!)^{-1/R} \geq 1/3$, $2^{-1-2\Sigma/R} \geq 2^{-3} = 1/8$, and, since $N \geq 1$, $N^{1-\Sigma/R} \geq N^{1-12/\sqrt{r}}$. $\square$

20.7 Counting candidates in a general interval

As in the ordinary-power count of Section 19, sum the spans of consecutive windows of $r+1$ candidates. Every adjacent gap is counted at most r times.

Corollary 20.7 (Semigroup interval count). *For $r \geq 144$ and $1 \leq H \leq N$,*

$$\#(\mathcal{C} \cap [N, N+H]) \leq r + 24rHN^{-1+12/\sqrt{r}}. \quad (71)$$

Proof. If s is the number of candidates and $s \leq r$, there is nothing to prove. Otherwise, Corollary 20.6 gives a lower bound $L = N^{1-12/\sqrt{r}}/24$ for each span $m_{i+r} - m_i$. Thus

$$(s-r)L \leq rH,$$

which is equivalent to (71). $\square$

20.8 An independent dyadic logarithmic-square bound

Corollary 20.8 (Explicit $O((\log X)^2)$ dyadic bound). *For every real $X \geq 1$,*

$$\#(\mathcal{C} \cap [X, 2X]) \leq 10000 \max\{1, \log X\}^2. \quad (72)$$

Proof. Put $L = \max\{1, \log X\}$ and

$$r = \lceil 144L^2 \rceil.$$

Then $r \geq 144$, $\sqrt{r} \geq 12L$, and

$$X^{12/\sqrt{r}} = \exp\left(\frac{12 \log X}{\sqrt{r}}\right) \leq e.$$

Apply (71) with $N = H = X$:

$$\#(\mathcal{C} \cap [X, 2X]) \leq r + 24rX^{12/\sqrt{r}} \leq (1 + 24e)r.$$

Since $L \geq 1$, $r \leq 145L^2$, and $(1 + 24e)145 < 10000$. $\square$

Status: [\[paper proof in this report\]](#)

The bound is numerically coarse by design. Its value is conceptual independence: it comes from positivity and integrality of generalized Vandermonde determinants, not from the prime-factor valuation rank of candidates. Refining constants is much less important than finding an additional source of determinant lower bound.

Status: [\[kernel-verified Lean\]](#) and [\[paper proof in this report\]](#)

The hierarchy is now kernel-accepted in `MixedExponentSemigroup`, `SemigroupDeterminant`, `FallingFactorialBound`, and `SemigroupWeightBound`. The final span and $10^4(\log X)^2$ deductions are accepted in `SemigroupGapBound` conditional on its explicit `NewtonFactorization` argument. That predicate is a visible input, not an axiom; the unconditional analytic proof remains the paper proof in this section.

21 The min-plus ceiling: why term-wise p -adic amplification cannot close a determinant argument

Every interpolation-determinant argument in this report squeezes a nonzero integer determinant D between a lower and an upper bound. A recurring proposal — registered in the deficit analysis of Section 22 as the most natural way to supply the missing logarithm — is to manufacture the lower bound from p -adic content: choose the nodes and the column exponents so that every monomial in the Leibniz expansion of D is divisible by a high power of 2 and of 3, and conclude

$$v_p(D) \geq V_p^* := \min_{\sigma} \sum_i v_p(E_{i\sigma(i)}), \quad |D| \geq 2^{V_2^*} 3^{V_3^*}. \quad (73)$$

This section records why the proposal cannot work, and — as importantly — corrects the quantitative form in which the obstruction has been stated.

21.1 The threshold

Fix a column set $S = \{(a_j, b_j)\}_{j=0}^r$ of distinct mixed exponents, put $\lambda_j = a_j + b_j\vartheta$ and $W = \sum_j \lambda_j$, and let D be the semigroup determinant $\det(c_i^{a_j} c_i^{b_j})$ built on the candidates $c_0 < \dots < c_r$ of a single dyadic block $[2^T, 2^{T+1})$. Every entry satisfies $\log_2 E_{ij} \leq (T+1)\lambda_j$, so the Leibniz bound gives

$$\log_2 |D| \leq (T+1)W + \log_2(r+1)! = TW(1 + o(1)). \quad (74)$$

An amplification scheme yielding a guaranteed lower bound $|D| \geq 2^V$ therefore needs $V > (T+1)W + \log_2(r+1)!$ for a contradiction. We call the right-hand side of (74) the *bar*.

21.2 The ceiling

The obstruction is not quantitative at all. It is the observation that the $\{2, 3\}$ -part of an integer divides the integer.

Theorem 21.1 (Min-plus ceiling). *Let $n \geq 1$ be a natural number, let $V_2, V_3 \geq 0$ satisfy $2^{V_2} \mid n$ and $3^{V_3} \mid n$, and suppose $n \leq 2^B$. Then*

$$V_2 + V_3 \vartheta \leq B.$$

More generally, for any finite set S of primes, $\sum_{p \in S} v_p(n) \log p \leq \log n$.

Proof. 2^{V_2} and 3^{V_3} are coprime and both divide n , so $2^{V_2} 3^{V_3} \mid n$ and hence $2^{V_2} 3^{V_3} \leq n \leq 2^B$. Taking $\log_2$ gives the claim. For general S the same argument applies to $\prod_{p \in S} p^{v_p(n)}$, which divides n by pairwise coprimality. $\square$

Status: [\[kernel-verified Lean\]](#) — `minPlusContent_le_of_le` and `minPlusContent_le_of_natAbs_le` in `MinPlusCeiling`, with the general prime-set form `primeContent_le_log` and the divisibility input `prod_primePow_dvd`.

Corollary 21.2 (Term-wise amplification is inert). *Let $D \neq 0$ be the determinant of an interpolation system, let V_2^*, V_3^* be any certified min-plus lower bounds as in (73), and let B be any correct archimedean upper bound $\log_2 |D| \leq B$. Then $V_2^* + V_3^* \vartheta \leq B$.*

Consequently no choice of node system, no choice of column set, and no enlargement of the prime set can make a term-wise divisibility bound contradict a correct archimedean bound.

The corollary is unconditional and needs no equidistribution input, no averaging, and no property of the candidates beyond $D \in \mathbb{Z} \setminus \{0\}$. Its content is that term-wise divisibility contributes nothing beyond what $|D| \geq 1$ already gives: it is a lower bound for $|D|$ derived from $|D|$ itself.

21.3 A correction to the averaging form of the obstruction

The obstruction has previously been argued by averaging: the minimum over permutations is at most the mean over permutations, the mean of $\sum_i v_p(E_{i\sigma(i)})$ is $\sum_j (a_j \bar{\nu} + b_j \bar{\mu})$ with $\bar{\nu}, \bar{\mu}$ the row-averaged valuations, and the depths n_k equidistribute over $[0, T]$, so $\bar{\nu} = T/2$. That computation is correct, and its conclusion — that the ceiling is governed by the *same* block equidistribution that defeats the shrinking-target route of Section 23 — is a genuine structural insight worth keeping: at the structural places the two walls of this problem are one wall.

The numerical conclusion drawn from it, however, is too strong. Bounding the 2-adic place and the ϑ -weighted 3-adic place separately gives $TW/2$ *each*; their sum is TW , which is the bar, not half of it. With $M = 2^\beta$ and $A = 3^\beta$ and the normalization $s_M = 2^{v_2(M)} 3^{v_3(M)}$, $s_A = 2^{v_2(A)} 3^{v_3(A)}$ for the 3-smooth parts, the mean bound reads

$$V^* \leq \frac{T}{2} \left[W + \frac{1}{\beta} \sum_j (a_j \log_2 s_M + b_j \log_2 s_A) \right], \quad (75)$$

and the bracket is at most $2W$, with equality exactly when M and A are both 3-smooth. Theorem 12.5 excludes that case, so the mean bound is strictly below the bar; but the deficiency it provides is proportional to the roughness of the outputs and degenerates as the outputs approach smoothness.

The factor $1/2$ requires $3 \nmid M$ and $2 \nmid A$. The matched-depth theorem of Section 5 excludes only the pair $2 \mid M$ and $3 \mid A$; the crossed pair $3 \mid M$ and $2 \mid A$ is not excluded, and it is exactly the configuration in which the ceiling climbs.

21.4 Calibration: the exact ceilings

The mean bound is not the true ceiling, because the minimum is genuinely lower than the mean. The exact constants were obtained in the continuation report by an optimal-transport computation, and they replace every heuristic figure that has circulated for this obstruction.

Write $\mathfrak{G}$ for the Gini norm of a normalized prime-valuation slope vector, and c_p for the slope vector at p . The termwise all-prime divisor fraction of the balanced mixed-semigroup determinant is exactly $1 - \frac{3}{8} \sum_p \mathfrak{G}(c_p) + o(1)$, and the primitive condition $\min(v_2(M), v_3(A)) = 0$ forces $\sum_p \mathfrak{G}(c_p) \geq 8/15$.

System	Termwise-divisor ceiling
Ordinary input/output Vandermonde product	< 0.8710491
Mixed single block, arbitrary branch	$4/5$
Mixed single block, cross-coprime	$7/10$
Mixed single block, places 2 and 3 only, cross-coprime	$3/10$
Complete triangular row grid	$\pi/4 = 0.7853982\dots$
First tied fibers at 2 and 3	$o(1)$ of the block budget

Two figures that have been proposed for this obstruction are not valid in the stated generality. A structural ceiling of $1/2$ requires dropping the depths $k u v_2(M)$ and $k v v_3(A)$ simultaneously, which assumes $v_2(M) = v_3(A) = 0$; the primitive theorem of Section 5 supplies only $\min(v_2(M), v_3(A)) = 0$. An all-prime ceiling of $2/3$ does not follow from “minimum $\leq$ average”: for an external positive weight proportional to $X_1 + X_2$ on the standard triangle one has $\mathbb{E}Y = 2/3$ and $\Phi(1, 1) = 4/15$, so the minimum assignment is $4/5$ of half the mean, not $2/3$.

An independent numerical check, computing exact assignment minima for the block node system over the places 2 and 3 only, in the cross-coprime branch and with a square-box column family, gives a ratio settling near 0.33 across block heights $T = 400, 1200, 3600, 10000$, consistent with the exact continuum value $3/10$ for those places. The same computation shows what happens outside the balanced primitive normalization: with a degenerate column family such as the pure- b columns $(0, j)$, and with A within $\log_2 3$ of being a power of three, the ratio rises to 0.9964. That configuration violates none of the kernel-verified constraints, but it saturates *tautologically* — the entries are then nearly powers of three, their 3-adic content is nearly their size, and (73) degenerates. The same degeneration occurs on the a -side at $b = 0$, where $D = \prod_{i < i'} (c_{i'} - c_i)$ and $v_2(D) = \sum_{i < i'} \min(n_i, n_{i'})$ exactly. This is why the balanced and primitive normalizations are not cosmetic: they are what make the $4/5$ and $7/10$ ceilings meaningful.

Status: [\[paper proof in this report\]](#) for the exact ceilings, from the continuation report’s transport computation — they follow from the closed-form Gini identity, not from experiment. [\[software computation\]](#) for the independent $3/10$ cross-check and the degenerate-family observation: exact rational node data, exact Jonker–Volgenant assignment minima, double-precision ϑ . Neither is needed for Corollary 21.2, which is [\[kernel-verified Lean\]](#) and unconditional.

21.5 Where cancellation can live

The ceilings bound only the *tropical* divisor. They say nothing about the true valuation $v_p(D)$, which can be larger when several permutations attain the minimum. The following observation localizes that possibility exactly.

Proposition 21.3 (A unique minimizer forbids cancellation). *Let $D = \det(E_{ij})$ be a nonzero integer determinant and let p be a prime. If a unique permutation σ_0 minimizes $\sum_i v_p(E_{i\sigma(i)})$, then $v_p(D) = \sum_i v_p(E_{i\sigma_0(i)})$.*

Proof. Divide the Leibniz expansion by p^{τ_p} with τ_p the minimum. The term indexed by σ_0 becomes a p -adic unit and every other term stays divisible by p , so the sum is a unit modulo p . $\square$

Status: [\[kernel-verified Lean\]](#) — `det_dvd_and_not_dvd_of_unique_min` in `MinPlusCeiling`, stated in the divisibility form ‘ $q^\tau \mid \det$ and $q^{\tau+1} \nmid \det$ ’ so that no p -adic valuation machinery is needed.

The corresponding tied-minimum statement is now also exact at the algebraic level. Suppose each Leibniz product has been written

$$\prod_i E_{i,\sigma(i)} = p^\tau q_\sigma$$

with $q_\sigma \in \mathbb{Z}$. Then the determinant factors as

$$\det E = p^\tau \sum_\sigma \operatorname{sgn}(\sigma) q_\sigma, \quad (76)$$

and, for every $k \geq 0$,

$$p^{\tau+k} \mid \det E \iff p^k \mid \sum_\sigma \operatorname{sgn}(\sigma) q_\sigma. \quad (77)$$

If p is prime and S is the tied minimum fiber, while the quotients outside S are divisible by p , the first extra factor is characterized exactly by the initial form

$$p^{\tau+1} \mid \det E \iff \sum_{\sigma \in S} \operatorname{sgn}(\sigma) q_\sigma = 0 \quad \text{in } \mathbb{Z}/p\mathbb{Z}. \quad (78)$$

Thus higher-cost terms disappear modulo p and all first-layer cancellation is concentrated in the tied fiber.

Status: [\[kernel-verified Lean\]](#) — `TropicalInitialForm` proves the factorization and all-depth divisibility criterion as `det_eq_pow_mul_signedLeibnizQuotientSum` and `det_pow_add_dvd_iff_signedLeibnizQuotientSum_dvd`; its prime first-layer forms are `det_pow_succ_dvd_iff_signedLeibnizQuotientSum_eq_zero`, `signedLeibnizQuotientSum_mod_eq_tiedFiberSum`, and `det_pow_succ_dvd_iff_tiedFiberSum_eq_zero`.

This module is a purely algebraic prerequisite and diagnostic. It proves neither that the tied sum vanishes nor that it has a large higher p -adic valuation; in particular it supplies no cancellation lower bound and does not close Conjecture 2.1.

The abstract block factorization needed for the report’s first-fiber calculation is also now kernel-verified. If a block map b describes the surviving permutations and A_k is the square submatrix on the fiber $b^{-1}(k)$, then

$$\sum_{\sigma \text{ preserving } b} \text{sgn}(\sigma) \prod_i A_{\sigma(i),i} = \prod_k \det A_k. \quad (79)$$

The formula allows varying block sizes and a fixed preliminary row permutation, whose sign is the only extra factor. When, after reindexing, each A_k consists of consecutive powers of nodes $x_{k,0}, \dots, x_{k,s_k-1}$, it becomes the exact Vandermonde product

$$\prod_k \prod_{0 \leq i < j < s_k} (x_{k,j} - x_{k,i}). \quad (80)$$

Composing (79) with the tropical quotient criterion gives an exact first-extra-factor test: under the explicit hypotheses that every quotient outside the block-preserving set vanishes modulo p and that each block-preserving quotient equals the corresponding full Leibniz product of A , $p^{\tau+1} \mid \det E$ if and only if $\prod_k \det A_k$ vanishes modulo p ; in the consecutive-power specialization the right side is precisely (80) modulo p .

Status: [\[kernel-verified Lean\]](#) — BlockFiberFactorization proves `sum_blockPreserving_eq_prod_blockDet`, `sum_basePerm_mul_blockPreserving_eq_sign_mul_prod_blockDet`, `sum_blockPreserving_eq_prod_vandermonde_of_reindex`, and `sum_blockPreserving_consecutivePowers_eq_prod_vandermonde`. TropicalBlockApplication packages the composition as `det_pow_succ_dvd_iff_prod_blockDet_eq_zero` and `det_pow_succ_dvd_iff_prod_vandermonde_eq_zero`.

The rank-one Beatty scaffolding is now kernel-verified separately. For matrices

$$E_{ij} = p^{n_i c(b_j)} A_{ij}, \quad n_i = T - \lfloor i\beta \rfloor,$$

with strictly decreasing Beatty depths, sorted block labels b_j , and strictly increasing block values c , `BeattyTropicalFiber` proves that the minimum assignments are exactly the block-preserving permutations. Every other assignment gains at least $\lfloor \beta \rfloor$, its normalized quotient is divisible by $p^{\lfloor \beta \rfloor}$, and the determinant has the exact expansion

$$\det E = p^\tau \left(\prod_k \det A_k + p^{\lfloor \beta \rfloor} Z \right).$$

The headline declarations are `rankOneAssignmentCost_eq_iff_blockPreserving`, `beattyAssignmentCost_add_floor_le_of_not_blockPreserving`, `beattyPower_pow_floor_dvd_quotient_of_not_blockPreserving`, and `det_beattyPowerMatrix_eq_pow_mul_blockDet_add_pow_floor_mul`. Thus the minimizing fiber, the tropical gap, and the full-product quotient identity no longer remain paper-only hypotheses. The generic scaled consecutive-power wiring is now formalized below, and `GlobalMixedPrefixCascade` closes the formerly paper-level finite specialization: its sorted box prefix is the genuine global first- N mixed-semigroup prefix, its coordinate fibers are lower-closed, and it constructs the lexicographic and colexicographic fiber equivalences needed for the dyadic and triadic determinant identities. What remains globally is the valuation and cancellation of the residual low-excess sums. No leading-order upper bound follows from the exact expansion.

The complete assignment cascade can now be grouped more economically than the raw Leibniz sum. A *row-block assignment* records only which row set is sent to each repeated column weight, with the prescribed fiber cardinalities. For the rank-one power matrix, the whole Leibniz fiber over such an assignment a factors as one sign times a product m_a of square unit minors. If $c(a)$ is its structural assignment cost and τ is the minimum, then

$$\det E = p^\tau \sum_a p^{c(a)-\tau} m_a. \tag{81}$$

More sharply, for every cutoff K this is the exact finite-depth decomposition

$$\det E = p^\tau \left(\sum_{c(a)-\tau < K} p^{c(a)-\tau} m_a + p^K \sum_{c(a)-\tau \geq K} p^{c(a)-\tau-K} m_a \right). \tag{82}$$

Thus modulo p^K every high-excess row partition disappears, and the surviving coefficients are explicit products of unit minors rather than sums over individual permutations. In the Beatty specialization, every noncanonical row partition has excess at least $\lfloor \beta \rfloor$.

Status: [\[kernel-verified Lean\]](#) — `RowBlockLaplaceExpansion` proves the fiber factorization `sum_permutationRowBlockAssignmentFiber_eq_sign_mul_prod_blockDet`, the normalized all-layer identity `det_rankOnePowerMatrix_eq_minPow_mul_sum_assignmentExcess`, the exact cutoff formula `det_rankOnePowerMatrix_eq_minPow_mul_truncatedAssignmentSum`, and the Beatty gap `beattyRowBlockAssignment_excess_ge_floor`.

This is an exact support theorem for every finite cascade depth. It gives no upper bound for the number of surviving row partitions by itself, no upper bound for the valuations of the residual minor-product quotients, and no noncancellation statement for their signed sum. It therefore reorganizes the remaining wall but does not bound $v_p(\det E) - \tau$.

There is, however, now an exact entropy bound for the first of those quantities. For a row partition a , let $P_k(a)$ be the excess of the total block value assigned to the first $k + 1$ rows over the oppositely sorted assignment, and define its crossing energy by

$$\mathcal{E}(a) = \sum_{k=0}^{N-1} P_k(a).$$

Every $P_k(a)$ is nonnegative, the vector $(P_k(a))_k$ determines a when the block values are distinct, and summation by parts identifies $\mathcal{E}(a)$ with the excess for the unit depth staircase. For Beatty depths,

$$c(a) - \tau \geq \lfloor \beta \rfloor \mathcal{E}(a). \quad (83)$$

Consequently, for $N, K \geq 1$, if $q = \lfloor (K - 1)/\lfloor \beta \rfloor \rfloor$, stars and bars gives the explicit finite support upper bound

$$\#\{a : c(a) - \tau < K\} \leq \binom{N + q - 1}{q}. \quad (84)$$

For fixed normalized depth q this is polynomial in N ; for growing q it is an explicit entropy bound, not a cancellation estimate.

Status: [\[kernel-verified Lean\]](#) — `AssignmentExcessGeometry` proves the exact identity `rowBlockCrossingEnergy_eq_reverseFinRank_excess`, the Beatty lower bound `beattyFloor_mul_rowBlockCrossingEnergy_le_excess`, the stars-and-bars injection `card_rowBlockAssignment_crossingEnergy_lt_choose_of_pos`, and the specialized bound `card_beattyRowBlockAssignment_excess_lt_of_pos`.

This controls how many row partitions can occur in the cutoff sum. It supplies no upper bound for their minor-product valuations and no noncancellation theorem for their signed sum, so (84) is cascade-support control rather than a proof of the conjecture.

The fixed-divisor contribution can be threaded through this entire expansion, rather than applied only to the canonical block. For a block-scaled consecutive-power unit matrix, let s_k be the fixed size of block k , define

$$U_2(s) = \sum_{0 \leq j < s} (j + v_2(j!)), \quad U_3(s) = \sum_{0 \leq j < s} \left(\left\lfloor \frac{j}{2} \right\rfloor + v_3 \left(\left\lfloor \frac{j}{2} \right\rfloor! \right) \right),$$

and put

$$\mathcal{U}_2 = \sum_k U_2(s_k), \quad \mathcal{U}_3 = \sum_k U_3(s_k).$$

Every row-block assignment has the same column-fiber sizes, so its minor product is divisible by $2^{\mathcal{U}_2}$, respectively $3^{\mathcal{U}_3}$. Factoring this common divisor from the exact row-partition expansion gives

$$2^{\tau_2 + \mathcal{U}_2} \mid \det E_2, \quad 3^{\tau_3 + \mathcal{U}_3} \mid \det E_3. \quad (85)$$

The row scale is allowed to depend on both the target block and the row; this includes the mixed factors M^{ku} in the dyadic blocks and A^{kv} in the triadic blocks. Thus the gain is genuinely uniform across every tropical layer, not a statement only about the first initial form.

Status: [\[kernel-verified Lean\]](#) — `RowBlockUnitDivisibility` defines `fiberScaledConsecutivePowerMatrix`, proves the common-minor divisors `two_pow_rowBlockOrderingExponent_dvd_minorProduct_fiberScaledConsecutivePowers` and `three_pow_rowBlockOrderingExponent_dvd_minorProduct_fiberScaledConsecutivePowers`, and packages (85) as `two_pow_minCost_add_rowBlockOrderingExponent_dvd_det_rankOnePowerMatrix` and `three_pow_minCost_add_rowBlockOrderingExponent_dvd_det_rankOnePowerMatrix`. The exponent is exactly the sum of the block gains from `UnitOrderingGain` and is independent of assignment cost. Consequently it factors every cutoff sum in (82), but it cannot distinguish layers or upper-bound the valuation of the residual signed sum.

There is also a cascade-independent fixed divisor of the unit determinants. A monic triangular change of columns from $1, X, X^2, \dots$ to an integer polynomial basis preserves the full evaluation determinant. On positive odd nodes with consecutive local degrees, the dyadic p -ordering basis

$$\prod_{0 \leq t < j} (X - (2t + 1))$$

extracts the column weight

$$j + v_2(j!),$$

while on positive nodes prime to three with consecutive local degrees, the ordering $1, 2, 4, 5, 7, 8, \dots$ extracts

$$\left\lfloor \frac{j}{2} \right\rfloor + v_3 \left(\left\lfloor \frac{j}{2} \right\rfloor! \right).$$

The resulting prime powers divide the *whole* univariate Vandermonde determinant, independently of which tropical layer survives. At the univariate block level this extracts a universal lower-order divisor rather than estimating successive initial forms.

Status: [\[kernel-verified Lean\]](#) — `AllLayerUnitBasis` proves the generic column-content and determinant-one column-change lemmas `prod_columnDivisor_dvd_det`, `pow_sum_columnWeight_dvd_det_of_unimodular_mul`, and `pow_sum_dvd_det_vandermonde_of_monicBasis`. Its strengthened unit specializations are `two_pow_sum_degree_add_factorialVal_dvd_det_vandermonde` and `three_pow_sum_halfDegree_add_factorialVal_dvd_det_vandermonde`.

The companion module `UnitOrderingGain` proves exact digit-sum identities for the block weights and the finite aggregate bounds

$$\sum_k U_2(s_k) \leq S^3, \quad 2 \sum_k U_3(s_k) \leq S^3$$

whenever there are at most S blocks and every block has size at most S (`sum_twoUnitBlockGain_le_cube` and `two_mul_sum_threeUnitBlockGain_le_cube`). For the report’s $O(L)$ fibers of size $O(L)$, the remaining paper-level block count therefore gives $O(L^3)$, equivalently $O(N^{3/2})$, lower order than $H_T \asymp T^{5/2}$. In fact, for the triangular cutoff $u + \vartheta v \leq L$ and $N \sim L^2/(2\vartheta)$, elementary summation gives the paper-level leading terms for the total gains

$$\mathcal{U}_2 \sim \frac{2^{3/2}}{3\sqrt{\vartheta}} N^{3/2} = 0.7488834673 \dots N^{3/2}, \quad \mathcal{U}_3 \sim \sqrt{\frac{\vartheta}{8}} N^{3/2} = 0.4451070799 \dots N^{3/2}.$$

The finite divisors and cubic aggregation are kernel-verified; the triangular fiber asymptotics and displayed constants are not Lean theorems. The determinant-one fiberwise change and its additive combination with the tropical minimum are kernel-verified for the generic scaled block model by `RowBlockUnitDivisibility`; the concrete triangular index bridge, including column permutation and transfer back to the raw mixed integer matrix, is kernel-verified under explicit finite fiber certificates in `MixedDeterminantUnitBridge`. The module `GlobalMixedPrefixCascade` now constructs those certificates for the genuine global first- N mixed-semigroup enumeration and proves the exact all-layer and cutoff-truncated dyadic and triadic formulas directly. The resulting all-layer gain neither supplies nor upper-bounds the remaining leading-order cancellation excess.

Consequently the missing fifth, or more, can arise only at primes where the assignment problem ties, and the tie geometry is explicit: for $p \neq 2, 3$ columns tie along the rational lines orthogonal to the external valuation vector, and in the cross-coprime branch the 2-adic cost depends only on u and the 3-adic cost only on v , so every vertical or horizontal column fiber creates equal-cost permutations. The continuation report evaluates the first such tied layer. The kernel-verified finite cube budget, combined with the elementary paper-level count of the Beatty fiber sizes, gives $O(T^2)$ against $H_T \asymp T^{5/2}$, hence $o(1)$ of the budget; the sharper quoted two-logarithm estimate improves this to $O(T^{3/2}(\log T)^2)$.

21.6 What this leaves

Three statements survive every normalization.

First, term-wise amplification is retired unconditionally against the crude entry bound. Corollary 21.2 is a structural identity, not a numerical near-miss, and it holds for every prime set, every node

system and every column set.

Second, the quantitative ceilings above say how much of the archimedean budget automatic divisibility can supply in each concrete construction, and none reaches 1. Adding every prime in the fixed support does not close any determinant; it improves the diagnosis from a heuristic “about one half” to a proved $4/5$.

Third — and this is the redirection — a successful p -adic proof has to be a *cancellation* theorem, not a divisibility-by-every-term theorem. Proposition 21.3 says exactly where such a theorem could live: inside the equal-cost fibers, and it would have to repeat through many tropical layers or be combined with a sharper archimedean free energy than the crude bar. Quantifying the gain of the refined divided-difference bounds of Sections 19 and 20 over the crude bar is the concrete companion question, and it concerns the archimedean side alone.

Status: [paper proof in this report] for the reading of the obstruction; [kernel-verified Lean] for Theorem 21.1, Corollary 21.2 and Proposition 21.3. The structural observation that the ceiling and the shrinking-target obstruction are governed by the same block equidistribution is recorded as an interpretive claim, not a theorem.

22 Why the new hierarchy still stops short

The semigroup determinant is much stronger locally than the ordinary-power determinant: its repulsion exponent approaches 1. It is therefore important to identify exactly why it does not prove the conjecture.

22.1 How many nodes are needed to approach scale X ?

The asymptotic exponent in (68) is

$$\delta_r := 1 - \frac{\Sigma_r}{R_r} = 1 - \frac{c_\vartheta}{\sqrt{r}} + O(r^{-1}), \quad c_\vartheta = \frac{4}{3}\sqrt{2\vartheta} = 2.3739044262\dots$$

For nodes near X , the lower span has the shape

$$X^{\delta_r} = X \exp\left(-\frac{c_\vartheta \log X}{\sqrt{r}} + O\left(\frac{\log X}{r}\right)\right). \quad (86)$$

To make the exponential loss in (86) bounded independently of X , one needs

$$\sqrt{r} \asymp \log X, \quad \text{that is,} \quad r \asymp (\log X)^2. \quad (87)$$

This is precisely why Corollary 20.8 has a logarithmic square.

22.2 How many nodes a counterexample supplies

Conditional on a counterexample, the exact dyadic block count established in Section 5 gives

$$\#(\mathcal{C} \cap [X, 2X]) = \frac{\log_2 X}{\beta} + O(1) \asymp \log X. \quad (88)$$

Thus the actual conditional population is smaller than the determinant threshold (87) by one factor of $\log X$.

Taking $r \asymp \log X$ in (86) yields only

$$X \exp(-C\sqrt{\log X}),$$

which is much smaller than the average spacing $X/\log X$ among the $\asymp \log X$ conditional candidates. The inequality is therefore compatible with the exact monoid even at the level of expected spacing.

22.3 Broader multiplicative intervals do not help

The cumulative conditional count below X is $\asymp (\log X)^2$, so one might try to use all those points at once. But their nodes range from 1 to X , and the gap theorem depends on the smallest node. Restricting to $[X^\gamma, X]$ with fixed $0 < \gamma < 1$ does retain $\asymp (\log X)^2$ points, but its additive length is $\asymp X$, while the lower scale is at most $X^{\gamma+o(1)}$. Again there is no contradiction.

This is not a mere defect of coarse constants. It reflects a mismatch between additive node geometry and the logarithmic triangular lattice of exponents (n, k) in $m = 2^n M^k$.

22.4 The determinant needs one more source of size

The archimedean argument uses only

$$D \in \mathbb{Z}_{>0} \implies D \geq 1.$$

If the exact monoid coordinates forced

$$D \equiv 0 \pmod{Q} \quad \text{with} \quad \log Q \gg R \log X / \sqrt{r},$$

then the lower bound $D \geq Q$ could remove the one-logarithm deficit. Alternatively, one could seek a determinant with total exponent weight $\Sigma = O(r)$ rather than $\Sigma \asymp r^{3/2}$, but the two-dimensional semigroup counting law $L(T) \asymp T^2$ makes that impossible using only distinct nonnegative mixed monomials.

This gives a concrete design specification:

A successful determinant proof must add nonarchimedean divisibility, introduce lower-weight functions outside the mixed-monomial semigroup, or exploit the exact two-dimensional row grid so efficiently that $O(\log X)$ nodes suffice.

Status: [paper proof in this report]

The deficit analysis of this section is a paper argument about the limits of the method, not a formalized theorem. It does not weaken any statement proved above; it explains why the explicit $O((\log X)^2)$ bound of Corollary 20.8 cannot be pushed to the conditional $\asymp \log X$ population by tuning constants alone. The conjecture remains open.

23 Further equivalent reformulations

Several equivalent forms are already formalized in the repository and were catalogued in Section 4. The exact conditional structure of Section 5 also yields useful paper-level reformulations that expose different possible proof technologies. Each restatement below is an exact equivalence, not a heuristic analogy: a proof of any one of them is a proof of Conjecture 2.1.

23.1 Integer points on a transcendental lattice curve

The candidate formulation of Section 2 says that

$$\{(m, n) \in \mathbb{N}^2 : n = m^\vartheta\} = \{(2^k, 3^k) : k \in \mathbb{N}_0\}. \quad (89)$$

The curve $n = m^\vartheta$ is strictly convex and transcendental, and (89) asserts that its only lattice points are the trivial exponential ones. The arbitrary-node repulsion theorem of Section 19 can be read as a lower bound for the volumes of lattice simplices cut from this curve; the deficit analysis of Section 22 measures exactly how far the current bound is from excluding the nontrivial points.

23.2 Vanishing determinant of four logarithms

Writing $M = 2^x$ and $A = 3^x$, a solution is a positive integer pair with

$$\log M \log 3 - \log A \log 2 = 0, \quad M, A \in \mathbb{N}, \quad \implies \quad (M, A) = (2^k, 3^k). \quad (90)$$

The hypothesis of (90) is a vanishing 2×2 determinant of logarithms of positive integers. This is the most direct bridge to the four-exponentials circle of ideas and to algebraic independence of logarithms, and it is the form in which the kernel-verified four-exponentials reduction is stated.

23.3 Primitive localized-radical exclusion

The kernel-verified radical form of Section 6 is

$$\forall w > 1 \text{ odd and power-primitive}, \quad \forall d > 0, a \geq 0, \quad 3^a(w^\vartheta)^d \notin \mathbb{N}. \quad (91)$$

If (91) fails with least rational-power index d , then $X^d - c/3^a$ is the exact minimal polynomial of w^ϑ over $\mathbb{Q}$ for the corresponding rational c .

This formulation is well suited to Kummer-theoretic, norm, and ramification investigations, because all redundant proper powers have been removed by primitivity and the algebraic degree of the hypothetical radical is known exactly rather than merely bounded.

23.4 Exact growth dichotomy

Write

$$B(T) := \#(\mathcal{C} \cap [2^T, 2^{T+1}))$$

for the T th dyadic block count. The repository formalizes the cumulative subquadratic equivalence of Section 5; the exact block formula there, namely

$$B(T) = 1 + \left\lfloor \frac{T+1}{\beta} \right\rfloor \quad \text{under failure with least generator } \beta, \quad (92)$$

gives several immediate variants.

Proposition 23.1 (Block-growth reformulations). *Conjecture 2.1 is equivalent to each of the following:*

1. the dyadic block counts $B(T)$ are bounded;
2. $B(T) = o(T)$;
3. $\liminf_{T \rightarrow \infty} B(T)/T = 0$;
4. $B(T) = 1$ for every $T \in \mathbb{N}_0$.

Proof. If the conjecture holds, every block contains only 2^T , so all four properties hold. If it fails, (92) gives $B(T) = 1 + \lfloor (T+1)/\beta \rfloor$, so $B(T)/T \rightarrow 1/\beta > 0$ and none of the first three properties holds. The fourth is visibly equivalent to the candidate statement. $\square$

This is useful because a future analytic theorem need not classify candidates individually. Any uniform $o(T)$ bound for the number of candidates in the T th dyadic block would prove the conjecture.

The exactness of (92) permits a sharpening of item (4) that is, at first sight, dramatically weaker than the conjecture and yet still equivalent to it.

Proposition 23.2 (Infinitely many singleton blocks suffice). *Conjecture 2.1 is equivalent to the statement*

$$B(T) = 1 \quad \text{for infinitely many } T \in \mathbb{N}_0.$$

Proof. If the conjecture holds then $B(T) = 1$ for every T , so certainly for infinitely many. Conversely, suppose the conjecture fails, and let $\beta > 0$ be the irrational least generator of the exact solution monoid. By (92),

$$B(T) = 1 \iff \left\lfloor \frac{T+1}{\beta} \right\rfloor = 0 \iff T+1 < \beta,$$

which holds for at most the finitely many indices $T < \beta - 1$. Hence failure forces $B(T) = 1$ for only finitely many T , and the contrapositive gives the equivalence. $\square$

Remark 23.3. Proposition 23.2 cannot be weakened further to “ $B(T) = 1$ for some T ”. That statement is already known unconditionally and carries no information: the kernel-verified zero-free region $2^x < 4096$ gives $\beta \geq 12$, hence $B(T) = 1$ for $T = 0, \dots, 10$, and the interval computation of Section 25 pushes the same conclusion to $T \leq 19$ at the weaker [software computation] trust level. The gap between “for some T ” and “for infinitely many T ” is exactly the gap between a finite check and the conjecture, because none of the present results bounds β from above: a counterexample with a very large least generator would have $B(T) = 1$ for every $T < \beta - 1$, that is, on any range a fixed finite computation could ever reach. Only the infinitude of singleton blocks is decisive.

Status: [\[kernel-verified Lean\]](#) for Proposition 23.1: BlockGrowthReformulations proves the bounded, little- o , liminf-zero, and identically one equivalences. [\[paper proof in this report\]](#) for Proposition 23.2, whose short paper proof rests on the same kernel-verified exact dyadic block formula.

23.5 Compact synchronized S -integral orbit

The S -integral orbit of Section 5 yields another equivalence, this time one that isolates the synchronization of the two denominators.

Proposition 23.4 (Compact-orbit reformulation). *Conjecture 2.1 fails if and only if there exist integers $M, A > 1$ and an irrational $\beta > 0$ such that, for every $k \geq 1$, the point*

$$\left(\frac{M^k}{2^{\lfloor k\beta \rfloor}}, \frac{A^k}{3^{\lfloor k\beta \rfloor}} \right) \quad (93)$$

lies on the compact arc

$$\mathcal{A} = \{(2^t, 3^t) : 0 \leq t < 1\}.$$

Under these conditions the orbit is dense in $\mathcal{A}$.

Proof. Failure supplies the least generator β , with $M = 2^\beta$ and $A = 3^\beta$; then (93) is exactly $(2^{\{k\beta\}}, 3^{\{k\beta\}})$, and density follows from irrational rotation.

Conversely, arc membership gives $t_k \in [0, 1)$ such that

$$k \log_2 M - \lfloor k\beta \rfloor = t_k = k \log_3 A - \lfloor k\beta \rfloor.$$

Thus $\log_2 M = \log_3 A =: \gamma$, and $0 \leq k\gamma - \lfloor k\beta \rfloor < 1$ for every k . Hence $\lfloor k\gamma \rfloor = \lfloor k\beta \rfloor$, so $|k(\gamma - \beta)| < 1$ for every k ; therefore $\gamma = \beta$. Consequently $M = 2^\beta$ and $A = 3^\beta$ are integers while β is nonintegral, so the conjecture fails. The same identity now shows that the orbit equals the irrational-rotation orbit and is dense. $\square$

Status: [\[kernel-verified Lean\]](#) —

`not_alaogluErdosConjecture_iff_exists_compactOrbit` and
`alaogluErdosConjecture_iff_no_compactOrbit` in `CompactOrbit`.

This formulation isolates the synchronized-denominator feature that a specialized S -integral counting theorem would need to exploit: the two coordinates are S -integral for disjoint sets of primes, yet their denominator exponents are forced to agree exactly at every height.

23.6 Two-dimensional exponential grid

Under failure, every solution coordinate (n, k) gives

$$m_{n,k} = 2^n M^k, \quad m_{n,k}^\vartheta = 3^n A^k,$$

and for every mixed exponent $a + b\vartheta$ with $a, b \in \mathbb{N}_0$,

$$m_{n,k}^{a+b\vartheta} = (2^a 3^b)^n (M^a A^b)^k. \quad (94)$$

Thus the candidate evaluation matrix factors through a two-dimensional row lattice and a two-dimensional column lattice. The conjecture can be viewed as the assertion that this integer exponential grid cannot exist unless (M, A) is a common integral power of $(2, 3)$. This is the most natural starting point for a bespoke 2×2 interpolation determinant that uses more than the arbitrary-node geometry of Section 19, and it is the structure the semigroup hierarchy of Section 20 exploits only partially.

23.7 An operator formulation on the prime-logarithm space

All the reformulations above are Diophantine. There is also a purely linear-algebraic one, which recasts the problem as a statement about multiplication by ϑ acting on a single $\mathbb{Q}$ -vector space.

Let

$$W := \text{span}_{\mathbb{Q}}\{\log p : p \text{ prime}\} \subset \mathbb{R},$$

a $\mathbb{Q}$ -vector space of countably infinite dimension, with the $\log p$ as a basis by unique factorization. Multiplication by ϑ does not preserve W , so consider the largest part of W that it does return to W :

$$U := \{w \in W : \vartheta w \in W\}.$$

This is a $\mathbb{Q}$ -subspace, and it is never zero: $\vartheta \log 2 = \log 3 \in W$, so $\mathbb{Q} \log 2 \subseteq U$.

Proposition 23.5 (Operator form). $\dim_{\mathbb{Q}} U = \text{rank } K$, where $K = \{v \in \mathbb{Q}_{>0} : v^{\vartheta} \in \mathbb{Q}\}$ is the multiplicative group of Section 12. Consequently the rational-output form (22) — and hence Conjecture 2.1 — is equivalent to

$$\dim_{\mathbb{Q}}\{w \in W : \vartheta w \in W\} = 1.$$

Proof. An element of W is $w = \sum_p c_p \log p$ with rational c_p , almost all zero. Clearing a common denominator N gives $Nw = \log v$ for a positive rational v . Now $\vartheta w \in W$ says $\vartheta \log v = \sum_p d_p \log p$ with rational d_p ; clearing a denominator M turns this into $\vartheta \log(v^M) = \log u$ with $u \in \mathbb{Q}_{>0}$, that is $(v^M)^{\vartheta} = u$, i.e. $v^M \in K$. So U is exactly the $\mathbb{Q}$ -span of $\log K$ inside W , whence $\dim_{\mathbb{Q}} U = \text{rank } K$. The conjecture in its rational-output form says $K = \langle 2 \rangle$, of rank one. $\square$

Status: **[paper proof in this report]** — an elementary translation, but a genuinely different framing: it replaces a Diophantine assertion by the statement that the operator “multiply by ϑ ” carries only a single line of W back into W .

Two remarks on what this framing does and does not give.

It explains why finite-dimensional arguments cannot apply directly. If some nonzero finite-dimensional subspace $V \subseteq W$ were ϑ -stable, then ϑ would be an eigenvalue of the rational matrix representing multiplication by ϑ on V , hence algebraic, contradicting transcendence. So a proof cannot proceed by exhibiting an invariant finite-dimensional subspace: none exists, unconditionally. The subspace

U above is not ϑ -stable — $\vartheta U \subseteq W$ but generally $\vartheta U \not\subseteq U$ — and this failure of stability is precisely what blocks the eigenvalue argument.

The natural stable candidate is empty for the same reason. Setting $U_\infty = \{w \in W : \vartheta^k w \in W \text{ for all } k \geq 0\}$ does produce a ϑ -stable space, but it is either zero or infinite-dimensional, and infinite dimension is automatic whenever it is nonzero: if $0 \neq w \in U_\infty$, the powers $\vartheta^k w$ are $\mathbb{Q}$ -linearly independent, since a relation $\sum_k c_k \vartheta^k w = 0$ with $w \neq 0$ would make ϑ algebraic. So the operator picture reproduces the transcendence barrier exactly, rather than circumventing it. Its value is diagnostic: it shows the obstruction is not a missing estimate but the absence of any finite-dimensional invariant structure to work with.

24 A consolidated catalogue of equivalent formulations

The preceding sections restate Conjecture 2.1 many times over: as a candidate-base identity, as a localized radical exclusion, as a counting dichotomy, as a statement about fractional parts, and as a closure property of the solution monoid. Because these restatements were introduced where they were needed rather than where they belong together, this section collects them in one place.

Throughout, $\vartheta = \log_2 3$,

$$\mathcal{S} = \{x \in \mathbb{R} : 2^x, 3^x \in \mathbb{Z}\}, \quad \mathcal{C} = \{m \in \mathbb{N} : m > 0, m^\vartheta \in \mathbb{N}\},$$

which correspond under $x \mapsto 2^x$ by the exponent–candidate bijection of Section 2, and

$$B(T) := \#(\mathcal{C} \cap [2^T, 2^{T+1}))$$

is the T th dyadic block count.

24.1 How to read the catalogue

Every item (E1)–(E26) below is an *exact biconditional*: a proof of any one of them is a proof of Conjecture 2.1, and a disproof of any one of them is a counterexample. None of them is proved. The status tag on each item records the strength of the evidence for the *equivalence*, never for the conjecture: **[kernel-verified Lean]** means the biconditional itself has been accepted by the Lean kernel, **[paper proof in this report]** means it is a paper proof in this report, and **[classical/open background]** means the equivalence is obtained by combining a paper proof here with a classical transcendence theorem.

Two adjacent notions are deliberately excluded from the catalogue and treated separately in Section 24.8 below: *sufficient conditions*, which are one-way implications from hypotheses that are themselves unproved, and *unconditional partial results*, which settle a bounded region rather than the statement. Conflating the three is the most common way to overstate what the corpus establishes.

24.2 Cluster A: normalized restatements

These three items are the conjecture rewritten in the coordinates in which the rest of the corpus works. They carry no arithmetic content beyond that bijection, and their value is that each suggests a different toolkit.

- (E1) [\[kernel-verified Lean\]](#) *Candidate form.* $\mathcal{C} = \{2^n : n \in \mathbb{N}_0\}$; equivalently, the only positive integers m with $m^\vartheta \in \mathbb{N}$ are the powers of two (Section 2).
- (E2) [\[kernel-verified Lean\]](#) *Vanishing four-log determinant.* The only positive integer pairs (M, A) with $\log M \log 3 - \log A \log 2 = 0$ are $(M, A) = (2^n, 3^n)$, $n \in \mathbb{N}_0$ (Section 2); this is the 2×2 logarithmic determinant whose entries include two logarithms of unknown integers, and the form in which the four-exponentials reduction of Section 7 is stated.
- (E3) [\[kernel-verified Lean\]](#) *One-parameter subgroup intersection.* The real one-parameter subgroup $\Gamma = \{(2^t, 3^t) : t \in \mathbb{R}\}$ of the multiplicative torus meets $\mathbb{N}^2$ exactly in $\{(2^n, 3^n) : n \in \mathbb{N}_0\}$ (Section 2); a geometric restatement of (E2), with the difficulty that the direction vector $(\log 2, \log 3)$ is transcendental rather than algebraic.

Status: [\[kernel-verified Lean\]](#) — the exponent–candidate bijection is `exists_twoBaseNaturalCandidate_of_two_three_rpow_integer` together with `TwoBaseNaturalCandidate.integer_powers`, and the equivalence (E1) is

`alaogluErdosConjecture_iff_candidates_are_powers_of_two.`

Items (E2) and (E3) are the same content transported along that bijection and along $m \mapsto (m, m^\vartheta)$; they carry no separate Lean identifier.

24.3 Cluster B: localized radicals

This cluster is the most compressed form of the problem known: a statement quantified over all real x becomes a statement about one integer radical at a time, with all 2-adic and all but one 3-adic degree of freedom eliminated. It is also the cluster closest to established transcendence theory.

- (E4) [\[kernel-verified Lean\]](#) *Localized radical exclusion.* There is no odd integer $u > 1$ with $u^\vartheta \in \mathbb{Z}[1/3]$ (Section 6).
- (E5) [\[kernel-verified Lean\]](#) *Integral localized form.* There is no pair (u, a) with $u > 1$ odd, $a \in \mathbb{N}_0$ and $3^a u^\vartheta \in \mathbb{N}$; equivalently, no candidate has odd part greater than one (Section 6).
- (E6) [\[kernel-verified Lean\]](#) *Primitive localized radical exclusion.* There is no power-primitive odd $w > 1$ and $d \geq 1$ with $(w^\vartheta)^d \in \mathbb{Z}[1/3]$ (Section 6); primitivity removes every proper-power redundancy, and under failure the minimal polynomial of w^ϑ is known exactly.
- (E7) [\[kernel-verified Lean\]](#) *Symmetric base-swapped form.* With $\eta = \log_3 2 = 1/\vartheta$, there is no integer $v > 1$ with $3 \nmid v$ and $v^\eta \in \mathbb{Z}[1/2]$ (Appendix C); the two normalizations produce the same least exponent, so the reduction is not an artifact of privileging the base 2.

Status: [\[kernel-verified Lean\]](#) — (E4) is

`alaogluErdosConjecture_iff_odd_rpow_not_localized`

in `Localization`, restated as

`alaogluErdosConjecture_iff_oddLocalizedRadicalExclusion`

in `RadicalReformulation`; (E5) is the independently proved candidate-level variant

`alaogluErdosConjecture_iff_no_odd_multiple_candidate;`

(E6) is

`alaogluErdosConjecture_iff_primitiveOddLocalizedRadicalExclusion`

in `PrimitiveRadical`; and (E7) is

`alaogluErdosConjecture_iff_threeFreeLocalizedRadicalExclusion`

in `RadicalReformulation`. The equivalences are kernel-verified; the exclusions on their right-hand sides are the open conjecture.

24.4 Cluster C: counting and block growth

The counting cluster replaces an integrality statement about individual candidates by an upper bound on how many of them there can be. It is attractive because an analytic theorem need never identify a candidate, and unattractive because the target quantities are extremely far from the arithmetic conclusion; Section 22 measures that distance precisely.

- (E8) [\[kernel-verified Lean\]](#) *Subquadratic cumulative growth.* $\#(\mathcal{C} \cap [1, 2^T])$ is $o(T^2)$ as $T \rightarrow \infty$ (Section 5).
- (E9) [\[kernel-verified Lean\]](#) *Bounded blocks.* The dyadic block counts $B(T)$ are bounded (Section 23).
- (E10) [\[kernel-verified Lean\]](#) *Sublinear blocks.* $B(T) = o(T)$ (Section 23).
- (E11) [\[kernel-verified Lean\]](#) *Sublinear along a subsequence.* $\liminf_{T \rightarrow \infty} B(T)/T = 0$ (Section 23).
- (E12) [\[kernel-verified Lean\]](#) *Singleton blocks.* $B(T) = 1$ for every $T \in \mathbb{N}_0$ (Section 23); an immediate restatement of (E1) blockwise.
- (E13) [\[paper proof in this report\]](#) *Infinitely many singleton blocks.* $B(T) = 1$ for infinitely many $T \in \mathbb{N}_0$ (Section 23).

Remark 24.1. Item (E13) is the sharpest illustration of how exact the conditional structure theory is. It looks dramatically weaker than (E12), yet the exact block formula $B(T) = 1 + \lfloor (T+1)/\beta \rfloor$ under failure with least generator β forces $B(T) = 1$ to hold only for the finitely many $T < \beta - 1$. It cannot be weakened further to “ $B(T) = 1$ for some T ”: that statement is already known unconditionally and carries no information, since the kernel-verified zero-free region gives $\beta \geq 12$. The gap between “for some” and “for infinitely many” is exactly the gap between a finite check and the conjecture.

Status: [\[kernel-verified Lean\]](#) for (E8) —

`alaogluErdosConjecture_iff_candidate_count_below_two_pow_subquadratic.`

[\[kernel-verified Lean\]](#) for (E9)–(E12) —

`alaogluErdosConjecture_iff_dyadicBlock_bounded,`
`alaogluErdosConjecture_iff_dyadicBlock_isLittle0,`
`alaogluErdosConjecture_iff_dyadicBlock_liminf_zero,` and
`alaogluErdosConjecture_iff_dyadicBlock_eq_one` in `BlockGrowthReformulations.`
[\[paper proof in this report\]](#) for (E13), whose paper proof rests on the kernel-verified exact block formula.

24.5 Cluster D: orbits and fractional parts

Under failure the solution set is an irrational-rotation orbit in disguise. This cluster converts the conjecture into the assertion that no such orbit can be arithmetically synchronized, and into three gap statements on the unit interval.

- (E14) [\[kernel-verified Lean\]](#) *Compact synchronized S -integral orbit.* There are no integers $M, A > 1$ and irrational $\beta > 0$ such that $(M^k/2^{\lfloor k\beta \rfloor}, A^k/3^{\lfloor k\beta \rfloor})$ lies on the compact arc $\{(2^t, 3^t) : 0 \leq t < 1\}$ for every $k \geq 1$ (Section 23). The two coordinates are S -integral for disjoint prime sets, yet their denominator exponents must agree exactly at every height.
- (E15) [\[kernel-verified Lean\]](#) *Avoided interval.* Some nonempty open interval $(u, v) \subseteq (0, 1)$ contains no fractional part $\{y\}$ of a nonintegral solution y (Section 4).
- (E16) [\[kernel-verified Lean\]](#) *Gap at the top.* The fractional parts of nonintegral solutions are bounded away from 1 (Section 4).
- (E17) [\[kernel-verified Lean\]](#) *Gap at the bottom.* The fractional parts of nonintegral solutions are bounded away from 0 (Section 4).

Items (E15)–(E17) are unconditional biconditionals, and they say something quantitative about the elementary gap bounds of Section 4: those bounds place $\{x\}$ away from 0 and 1 by margins that decay like $2^{-\lfloor x \rfloor}$, and the equivalences show that *any* height-independent margin, however small and wherever located, is already the full conjecture. The decay is therefore essential rather than an artifact of the estimate.

Status: [\[kernel-verified Lean\]](#) for (E14) —

`alaogluErdosConjecture_iff_no_compactOrbit` in `CompactOrbit.` [\[kernel-verified Lean\]](#) for (E15)–(E17) —

`alaogluErdosConjecture_iff_fract_avoids_interval,`

`alaogluErdosConjecture_iff_fract_bounded_away_one,`

`alaogluErdosConjecture_iff_fract_bounded_away_zero`

in `FractionalPartDensity.`

24.6 Cluster E: closure, smoothness, and algebraicity

The last cluster is the least classical and, for that reason, the most interesting as a source of new attacks. It says that the conjecture is equivalent to the failure of a single nonlinear closure property: the solution monoid, which is trivially closed under addition, must not also be closed under multiplication.

- (E18) **[paper proof in this report]** *Multiplicative closure.* $\mathcal{S}$ is closed under multiplication.
- (E19) **[paper proof in this report]** *Square closure.* $\mathcal{S}$ is closed under squaring.
- (E20) **[paper proof in this report]** *Star closure.* $\mathcal{C}$ is closed under the operation $m \star n = 2^{\log_2 m \log_2 n}$.
- (E21) **[paper proof in this report]** *Star idempotence.* Every $m \in \mathcal{C}$ satisfies $m \star m \in \mathcal{C}$.
- (E22) **[paper proof in this report]** *Smooth outputs.* Every $x \in \mathcal{S}$ has both 2^x and 3^x 3-smooth.
- (E23) **[paper proof in this report]** *Smooth product output.* Every $x \in \mathcal{S}$ has 6^x 3-smooth.
- (E24) **[classical/open background]** *Square-exponent algebraicity.* Every $x \in \mathcal{S}$ has $6^{x^2} \in \overline{\mathbb{Q}}$.
- (E25) **[classical/open background]** *Split square-exponent algebraicity.* Every $x \in \mathcal{S}$ has both $2^{x^2} \in \overline{\mathbb{Q}}$ and $3^{x^2} \in \overline{\mathbb{Q}}$.
- (E26) **[classical/open background]** *Mixed-product algebraicity.* All $x, y \in \mathcal{S}$ satisfy $6^{xy} \in \overline{\mathbb{Q}}$.

Remark 24.2 (Manufacturing a third base). Items (E22)–(E26) admit a concrete reading. By the kernel-verified 3-smooth classification of Section 4, a hypothetical nonintegral solution β admits an integer base B with $B^\beta \in \mathbb{Z}$ only when B is 3-smooth. It would therefore suffice to manufacture one non-3-smooth positive integer B with $B^\beta \in \mathbb{Z}$. The obvious candidates $B = 2^\beta$, $B = 3^\beta$, $B = 6^\beta$ have β -th powers 2^{β^2} , 3^{β^2} , 6^{β^2} , so “manufacture a third base” is exactly the square-closure problem of (E19) in different language. Whenever the chosen base is not 3-smooth its β -th power is already known to be transcendental by six exponentials, so any proof of algebraicity would finish the conjecture at once. What is missing is a single bridge:

derive one nonlinear algebraic value from the two linear integer values.

Neither unique factorization, nor the ordinary exponent laws, nor linear forms in logarithms supplies it.

Status: **[paper proof in this report]** for (E18)–(E23) — paper proofs in this report, from the closure equivalences and the output-smoothness theorem; no Lean formalization yet. **[classical/open background]** for (E24)–(E26) — the equivalences combine those paper proofs with the classical six exponentials theorem, which supplies transcendence of a non-3-smooth base raised to the hypothetical exponent. None of the nine items has been accepted by the kernel.

24.7 Which formulations are worth attacking

The twenty-six items are not equally promising, and it is worth saying so plainly. Cluster C is exact but analytically distant: the deficit analysis of Section 22 shows the determinant hierarchy

reaching $O((\log X)^2)$ while a counterexample supplies $\asymp \log X$ candidates, a full logarithm short of the $o(T)$ that (E10) would need. Cluster A is a change of coordinates, useful for orientation and for locating the four-exponentials barrier, but not itself a method. Cluster B is closest to established transcendence theory and is where a Kummer-theoretic, norm, or ramification argument would land; its smallest open instance is the irrationality of $3^{\log_2 3}$. Cluster D isolates a synchronization phenomenon that a bespoke S -integral counting theorem might exploit. Cluster E is the only one whose hypotheses are elementary and nonlinear at the same time, and it is the natural home of any attempt to manufacture a third base.

24.8 Equivalences versus sufficient conditions

An equivalence and a sufficient condition play very different roles, and the distinction is easy to lose in a corpus this size. Items (E1)–(E26) are biconditionals: each is exactly as strong as Conjecture 2.1, so proving any of them proves the conjecture and no weakening of any of them is safe without checking that the reverse implication survives. The statements below are strictly one-way. Each is a hypothesis that is itself open, each implies the conjecture, and none is known to follow from it — so none belongs in the catalogue.

- (S1) *The real four exponentials conjecture.* If for all $\mathbb{Q}$ -linearly independent x_1, x_2 and $\mathbb{Q}$ -linearly independent y_1, y_2 some $e^{x_i y_j}$ is transcendental, then Conjecture 2.1 holds. The implication is kernel-verified (Section 7); the hypothesis is open, and is the reason no routine combination of present methods closes the problem.
- (S2) *Schanuel’s conjecture.* It implies algebraic independence of the logarithms of the primes, hence (S3), hence Conjecture 2.1 (Section 12).
- (S3) *Prime-log-product independence.* If the real numbers $\log p \log q$, indexed by pairs of primes $p \leq q$, are linearly independent over $\mathbb{Q}$, then Conjecture 2.1 holds (Section 12). This criterion mentions no exponentials at all, is strictly weaker than Schanuel, and is not known to imply or to follow from (S1).

Status: [\[kernel-verified Lean\]](#) for (S1) —

`alaogluErdosConjecture_of_realFourExponentialsConjecture`

in `FourExponentialsReduction`, which formalizes the implication without assuming its hypothesis. [\[paper proof in this report\]](#) for the Schanuel implication in (S2).

[\[kernel-verified Lean\]](#) for the conditional implication in (S3): `LogProductIndependence` records the independence hypothesis as an explicit `Prop`, never asserted, with headline statement

`alaogluErdosConjecture_of_primeLogProductIndependence.`

The finite prime-support form of (S3) is weaker still and correspondingly more approachable: if the products $\log p \log q$ over a fixed finite prime set $S \ni 2, 3$ are $\mathbb{Q}$ -linearly independent, then no counterexample has both outputs supported on S . For $S = \{2, 3\}$ this reduces to $\mathbb{Q}$ -linear independence of $1, \vartheta$ and ϑ^2 , which would exclude every 3-smooth counterexample. That is a partial exclusion, not an equivalence.

A third category deserves the same care. Unconditional partial results settle a bounded region and are neither equivalences nor sufficient conditions: the kernel-verified zero-free region gives $2^x \geq 4096$, hence $x \geq 12$, for every nonintegral solution (Section 25), [\[kernel-verified Lean\]](#); the directed-rounding scan through 2^{27} extends the same conclusion at the weaker trust level of a software computation, [\[software computation\]](#), and is not a kernel proof. The kernel-verified exponent bound remains $x \geq 12$. No item in this section, and no combination of them, proves Conjecture 2.1; the catalogue records where a proof could enter, not that one has.

25 Independent finite verification

Exactly one zero-free region in this report is a kernel certificate. Lean proves that a solution with $2^x < 4096$ forces $x \in \mathbb{Z}$, so that a nonintegral solution satisfies $x \geq 12$; the statement lives in `FiniteCheck4096` as

`twelve_le_of_not_integer_of_two_three_rpow_integer`

and it is the only finite exclusion that any other statement here may use without adding trust assumptions. [\[kernel-verified Lean\]](#)

One further finite computation pushes the excluded range far past 4096: a directed-rounding exclusion through 2^{27} , described in Section 40. It supersedes the earlier interval scans, which are not restated here. It is a software result: it does not replay in the Lean kernel, it is not interchangeable with `FiniteCheck4096`, and it does not raise the kernel-verified exponent bound above $x \geq 12$. They are recorded here because they strengthen the explicit conditional picture, test the conditional models of Section 5 against data, and supply the extremal cases that a future generated certificate must clear.

25.1 What each scan must certify

Write $\vartheta = \log_2 3$, and let m be a positive integer that is not a power of two. Excluding m means proving $m^\vartheta \notin \mathbb{N}$, and it suffices to exhibit an integer n with $n < m^\vartheta < n + 1$. Exponentiation is never performed: the pair of strict inequalities is equivalent, after taking logarithms and clearing the positive factor $\log 2$, to

$$\log m \log 3 - \log n \log 2 > 0, \tag{95}$$

$$\log(n + 1) \log 2 - \log m \log 3 > 0. \tag{96}$$

Every logarithm and every product occurring in the scanned ranges is positive, so the required rounding directions are monotone and unambiguous.

In both scans an ordinary floating-point evaluation merely *proposes* $n \approx \lfloor m^\vartheta \rfloor$. Once certified lower bounds for (95) and (96) are both strictly positive, the floating proposal has no remaining logical role: it is a witness generator, not a step of the argument. A locator error can therefore create a spurious failure report, but it cannot create a spurious exclusion. Powers of two are skipped rather than tested, since 2^q has the integral value $(2^q)^\vartheta = 3^q$; an exactly integral value at a non-power of two would satisfy neither strict inequality for any n and would be reported as a failure.

25.2 Superseded scans

Two earlier software exclusions — an interval-arithmetic scan through 2^{20} and a directed-rounding scan through 2^{26} — have been absorbed by the extension through 2^{27} of Section 40, which uses the same directed-rounding method, the same margin conventions and the same manifest discipline over the wider range. Only the strongest exclusion is stated in this report; the superseded ranges and their chunk tables are not reproduced.

25.3 The trust boundary, stated plainly

Kernel. $2^x < 4096$ and $2^x, 3^x \in \mathbb{Z}$ imply $x \in \mathbb{Z}$; a nonintegral solution has $x \geq 12$. [[kernel-verified Lean](#)] `FiniteCheck4096`.

Interval scan. Under CPython and `mpmath` interval arithmetic, a nonintegral solution has $x > 20$. [[software computation](#)]

Directed-rounding scan. Under the C source, GCC code generation, OpenMP scheduling, the MPFR and GMP shared libraries, the x86-64 arithmetic environment, and the operating system, a nonintegral solution has $x > 27$. [[software computation](#)]

MPFR’s correct-rounding contract makes the numerical reasoning crisp, and the interval scan is reproducible from a short script, but neither removes the software assumptions listed above. Nothing in Section 5 or Section 27 that is tagged [[kernel-verified Lean](#)] depends on either scan; wherever a numeric threshold is used inside a kernel-verified statement, that threshold is 12.

A drafted extension of the kernel region. An attempt to raise the kernel-checked region from $2^x < 4096$ to $2^x < 16384$ exists in draft. It follows the `FiniteCheck4096` pattern, splitting the enlarged table into four chunks so that no single `decide` call carries the whole certificate. The certificate rows themselves verify: each chunk’s Boolean recursion evaluates and each row’s exact integer comparison is accepted. What has not been accepted is the assembly — the per-chunk length lemmas that glue the four tables into a single covering statement exceeded the elaborator’s recursion limit, so the module does not compile end to end and no theorem is available for use. Until that is resolved the kernel bound is unchanged at $x \geq 12$. [[Lean draft, not yet accepted](#)]

From scan to kernel certificate. The repository’s own certificate architecture (Section 4) shows the stronger trust model: generate rational certificates externally, revalidate them independently, and replay only exact integer comparisons in Lean by `decide`. The next step is certificate extraction, not a larger blind scan. For each block one records rational lower and upper approximants to ϑ together with integer inequalities sufficient to replay every exclusion, exploiting monotonicity and convexity so that a single certificate covers a run of consecutive m rather than storing one record per input. A scalable formal certificate should share convergent enclosures across ranges, chunk the data modules to bound elaboration memory — the failure mode of the 16384 draft is precisely that the chunk bookkeeping, not the arithmetic, is the expensive part — and rely on verified binary powering rather than giant literals. The extremal pairs identified by both scans predict the required tier depth. A practical intermediate target is a generated, kernel-replayed certificate beyond 4096 (with the existing 16384 draft as an engineering test), followed by a hierarchical certificate that approaches the current software frontier 2^{27} .

26 Attempts at a full proof and their exact failure points

This section records the principal avenues that were pushed until they broke, together with the exact place where each one breaks. The purpose is not merely negative. Every failure point names a quantitative or structural feature that a successful proof must improve, and several of them turn out to be the same feature seen from different sides; the last subsection gives that common feature its exact name.

26.1 Collapsing rank two to rank one

The multiplicative-rank bound is kernel-verified: the exponent vectors of the candidate set span a subgroup of rank at most two. Under failure the rank-two possibility is not vague but known exactly, $\mathcal{C} = \langle 2, M \rangle$, with M the primitive generator of Section 5. Nothing in the rank machinery distinguishes this free rank-two monoid from the allowed monoid of powers of two.

Failure point. A rank bound is an upper bound on dimension; it supplies no mechanism for forcing the dimension down to one. Excluding the second generator *is* the original transcendence problem, restated.

26.2 Rational approximation to a least generator

Assume failure and let β be the least nonintegral solution, with $M = 2^\beta$ and $A = 3^\beta$. Let p/q approximate β and put $\varepsilon = q\beta - p$. Then

$$M^q - 2^p = 2^p(e^{\varepsilon \log 2} - 1), \quad (97)$$

$$A^q - 3^p = 3^p(e^{\varepsilon \log 3} - 1). \quad (98)$$

Both left-hand sides are integers, so a nonzero value has absolute value at least 1, which for small ε gives bounds of the shape

$$|\varepsilon| \gtrsim 2^{-p} \quad \text{and} \quad |\varepsilon| \gtrsim 3^{-p}.$$

These are exponentially tiny in q . Continued fractions only guarantee infinitely many approximants with $|\varepsilon| \ll q^{-1}$, and Baker-type lower bounds for ordinary linear forms in logarithms are polynomial in the relevant heights [3, 10]. Neither scale conflicts with a lower bound as weak as M^{-q} or A^{-q} . Multiplying (97) by (98), or eliminating ε between them, only recovers $\log_2 3$ to first order; the higher terms are far too small to create an integer contradiction.

Failure point. The two integer gaps are correlated rather than independent, so the two exponential inequalities carry essentially one piece of information, not two.

Observation 26.1. Any rational-approximation proof needs more than “a nonzero integer has absolute value at least one.” It needs a divisibility lower bound growing exponentially with q , or an approximation to β exponentially better than continued fractions supply in general.

26.3 Equidistribution against exponentially shrinking targets

At paper level the orbit $\{k\beta\}$ is equidistributed modulo 1 (Section 5); the finite-Fourier convergence needed for Weyl’s criterion, covering every fixed mode and every finite trigonometric polynomial, is

kernel-verified. Meanwhile, integrality of 2^{n+t} between adjacent integral powers pushes the fractional part t away from both endpoints by an amount comparable with 2^{-n} , and the base-3 condition imposes a comparable 3^{-n} restriction; at height k the admissible target width is of exponential scale M^{-k} , a bound also kernel-verified in the shrinking-target module.

Equidistribution controls visits to fixed intervals and, with effective discrepancy, to polynomially shrinking ones along subsequences. It does not force visits to a summable family of targets of length e^{-ck} : the Borel–Cantelli heuristic predicts only finitely many such visits for a typical rotation, and numbers with bounded partial quotients show that equidistributed multiples can avoid exponentially small targets forever. The exact conditional monoid is perfectly compatible with that behaviour.

Failure point. The blockwise Weyl theorems are structurally correct but point the wrong way: they describe how already existing solutions distribute inside each block, not how likely arbitrary exponents are to hit the integer lattice. Closing the route needs Diophantine information specific to the transcendental β with $2^\beta, 3^\beta \in \mathbb{N}$, not generic metric theory.

26.4 Finite differences and additive progressions

For $f(t) = t^\vartheta$ the repository proves the exact mean-value formula for every forward difference order r ,

$$\Delta_h^r f(n) = (\vartheta)_r h^r \xi^{\vartheta-r} \quad (n < \xi < n + rh),$$

so that if all $r+1$ values $f(n), f(n+h), \dots, f(n+rh)$ are integers, the left side is an integer; under $|(\vartheta)_r| h^r < n^{r-\vartheta}$ its absolute value lies strictly between 0 and 1. This is a sharp uniform obstruction to candidate arithmetic progressions, and its rational effective three-term shadow $h^{400} \leq n^{83}$ is kernel-verified together with the arbitrary-order form.

The difficulty is combinatorial rather than analytic. Conditional failure produces only $O(\log X)$ candidates per multiplicative block, far below the positive-density regime of Roth-type theorems, and a rank-two multiplicative monoid need not contain long additive progressions at all: many points in a short interval need not be equally spaced. A three-term progression of candidates would solve the S -unit equation

$$2^{n_1} M^{k_1} + 2^{n_3} M^{k_3} = 2^{n_2+1} M^{k_2},$$

and general S -unit theory yields finiteness statements, not the required nonexistence.

Failure point. Sparse exact structure and sharp local curvature are mutually consistent. The arbitrary-node determinant of Section 19 removes the equal-spacing hypothesis, but its packing exponent is still too weak to contradict the exact conditional count.

26.5 Radical fields and conjugates

The normalized radical $E = w^\vartheta$ is algebraic with exact binomial minimal polynomial $X^d - c/3^a$, and its conjugates are $\zeta_d^j E$; degree, norm, and denominator support are all kernel-verified (Section 6). What this does not supply is a way to transport the real analytic identity $\log E = \vartheta \log w$ to the other conjugates.

Failure point. The chosen positive real logarithm is not Galois-equivariant: the remaining roots differ by complex logarithmic periods, and taking norms merely reproduces $E^d = c/3^a$. Ramification is equally inconclusive, since the equation relates the constrained primes through a quotient of

logarithms rather than through a field embedding; standard Kummer theory sees the radical equation but not the special value $\vartheta = \log_2 3$.

Observation 26.2. A radical-field proof needs a genuinely mixed invariant: something simultaneously sensitive to the algebraic conjugates of E and to the logarithmic relation among $2, 3, w, E$. At present the algebraic and the real-analytic arguments meet only at the four-logarithm determinant.

26.6 The compact S -integral orbit

From the least generator define

$$U_k := \frac{M^k}{2^{[k\beta]}} = 2^{\{k\beta\}} \in \mathbb{Z}[1/2], \quad V_k := \frac{A^k}{3^{[k\beta]}} = 3^{\{k\beta\}} \in \mathbb{Z}[1/3].$$

The points (U_k, V_k) lie on the compact analytic arc $\mathcal{A} = \{(2^t, 3^t) : 0 \leq t < 1\}$ and are dense there because β is irrational. This is the compact-orbit reformulation of Section 23, and it invites $\mathfrak{o}$ -minimal and S -integral point counting; Butler relates special cases of Wilkie-type counting to real three- and four-exponentials statements [5].

Failure point. The counting is not strong enough by an enormous margin. The logarithmic heights of the orbit points grow linearly in k , so the arc carries only about $\log H$ orbit points of height at most H , far below the subpolynomial count H^ε that Pila–Wilkie style theorems already tolerate outside algebraic parts. What would suffice is a denominator-sensitive refinement counting only points whose coordinate denominators are supported exactly on $\{2\}$ and on $\{3\}$ with synchronized exponents; no theorem of that strength appears in the inspected corpus.

26.7 Forcing an auxiliary base

If one could exhibit an integer $b > 1$ multiplicatively independent of 2 and 3 with b^x algebraic — or rational with suitably controlled denominator — then the six exponentials theorem, or the repository’s rational-third-base theorem, would force $x \in \mathbb{Q}$ and hence $x \in \mathbb{Z}$. The five exponentials theorem [10, 8] similarly makes $e^{\gamma x}$ transcendental for every nonzero algebraic γ , given the four algebraic corner exponentials, and the kernel-verified 3-smooth classification is the formalized integer-level shadow of the six exponentials theorem.

Failure point. The hypotheses provide no such b . The natural auxiliary quantities $p^x = e^{x \log p}$ have transcendental coefficient $\log p$, so the five exponentials theorem does not convert a divisor of M or A into a contradiction; and the six exponentials theorem controls b^x , not the distinct question of whether b divides the integer 2^x . The obvious constructions $b = 2^u 3^v$, and any construction from M and A , collapse to the same rank-two logarithmic span. An auxiliary base must come from an arithmetic operation whose exponential behaviour stays controlled while its logarithm leaves that span — an algebraic independence problem in disguise.

26.8 Exploiting the primitive output pair

The primitive pair (M, A) is canonical and rigid: kernel-verified theorems show that it admits no matched depth, no common perfect power, and no nontrivial affine decomposition, and that the odd core and its coordinates are unique. Every structural handle one might hope to grip has therefore been shown to be absent by design rather than by accident.

Failure point. After all of that normalization, what remains is exactly the single relation $\log M \log 3 = \log A \log 2$, a product of logarithms lying outside the scope of linear-forms machinery.

26.9 The linear-form-in-log-products formulation of the wall

The obstruction can be given one further exact name. A counterexample is a pair of integers $m, A \geq 2$ with $\log m \log 3 = \log A \log 2$. Expanding both integers over their prime factorizations $m = \prod_p p^{m_p}$ and $A = \prod_p p^{a_p}$, failure of Conjecture 2.1 is equivalent to the nontrivial vanishing of

$$D = \sum_p (m_p \log 3 - a_p \log 2) \log p, \quad (99)$$

an *integer-coefficient linear form in the products of two prime logarithms* $\{\log p \cdot \log 2, \log p \cdot \log 3\}$. Baker’s theory bounds linear forms in logarithms with *algebraic* coefficients [3]; here every coefficient of $\log p$ is itself a transcendental logarithm, and no lower-bound technology exists for such forms. Indeed even the $\mathbb{Q}$ -linear independence of the family $\{\log p \cdot \log q : p < q \text{ prime}\}$ is open — it follows from, and is essentially a fragment of, the four exponentials circle of conjectures. Section 12 turns exactly that independence statement into a sufficient condition for the conjecture.

Failure point. Every route examined above — determinants, product formulas in the radical field, auxiliary exponentials, equidistribution against shrinking targets, orbit counting — terminates at a special case of this one independence statement. A genuine proof therefore requires either a fundamentally new lower bound for linear forms in log-products, or a mechanism exploiting the positivity and integrality of the coefficients (m_p, a_p) in (99) that has no analogue in the general four-exponentials setting.

27 Consolidated research program

The two source programs overlapped substantially; the exact structure theory of Section 5 settles some directions outright and re-prioritizes the rest. The list below is ordered by how directly each item attacks the quantitative deficit of Section 22 — the one missing logarithm between what the determinant hierarchy delivers and what a contradiction requires.

27.1 Priority A: p -adic amplification of semigroup determinants

The semigroup determinant $D = \det(m_i^{a_j} (m_i^\vartheta)^{b_j})$ is a positive integer, and under the exact monoid every entry has explicit valuations obtained from $m_i = 2^{n_i} M^{k_i}$ and $m_i^\vartheta = 3^{n_i} A^{k_i}$:

$$\begin{aligned} v_2(\text{entry}_{ij}) &= a_j(n_i + v_2(M) k_i) + b_j v_2(A) k_i, \\ v_3(\text{entry}_{ij}) &= a_j v_3(M) k_i + b_j(n_i + v_3(A) k_i). \end{aligned}$$

What this would buy is decisive: archimedean reasoning offers only the universal lower bound $|D| \geq 1$, whereas a systematic p -adic factor supplies $|D| \geq p^L$ and could erase the missing logarithm outright.

Question 27.1. *Can one choose $O(\log X)$ candidate rows and mixed-monomial columns so that the resulting nonzero determinant satisfies $\log |D|_p^{-1} \gg (\log X)^2$ for $p = 2$ or $p = 3$, while the archimedean estimate remains $\log |D|_\infty = o((\log X)^2)$?*

The difficulty is that a naive row or column factor is lost because the exponent set contains $(0, 0)$; only more subtle min-plus determinant estimates, applied after suitable row and column finite-difference operations, have a chance. The exact affine primitivity and the gcd condition on the normalized degrees should be fed into the estimate rather than discarded.

27.2 Priority B: exploit the full (n, k) grid, not sorted magnitudes

The arbitrary-node argument forgets the canonical coordinates and retains only the sorted integers m_i . Restoring them exposes rectangular, often Kronecker-like, bivariate exponential matrices, which is exactly the balanced 2×2 setting where ordinary six exponentials estimates stop; integrality, positivity, and the one-depth-zero normalization may create a margin absent from the general four-exponentials problem. The right experiment is not another large generic determinant but an optimized minor selected from triangular coordinate regions $n + k\beta \leq T$ and $a + b\vartheta \leq L$, with exact accounting of archimedean size, zero estimates, and 2- and 3-adic valuations. The hard part is that the optimization is genuinely two-parameter and the useful shapes are not predicted by any current heuristic, so a search must precede a theorem.

The continuation in Section 65 turns this priority into an exact fewnomial target. On R distinct counterexample points, every vanishing auxiliary polynomial has at least $R + 1$ monomials; therefore one must either construct a relation with one fewer monomial than ordinary interpolation would use, or exploit additional arithmetic structure not measured by support. On the $2N$ grid, the BK-degree calibration shows a factor-four gap between the extremal interpolant and the general Matrix Coefficient target. The boundary-dilation defects are the most concrete finite compatibility equations currently available for attacking that gap.

27.3 Priority C: denominator-sensitive counting on the compact arc

The compact-orbit reformulation asks for infinitely many points on $y = x^\vartheta$ in $[1, 2) \times [1, 3)$ whose first denominator is a power of 2, whose second denominator is a power of 3, and whose denominator exponents are synchronized by the same integer $\lfloor k\beta \rfloor$. A theorem asserting that a compact transcendental arc with nonzero curvature carries $o(H)$ points of logarithmic height at most H with denominators supported on two prescribed disjoint prime sets and a common exponent parameter would close the case, since the failed-conjecture orbit has $\asymp H$ such points. What makes it hard is that standard rational-point counting is sensitive to denominator *size* but not to denominator *support* or synchronization; Butler’s connection between Wilkie-type statements and real exponential conjectures [5] suggests this direction sits on the transcendence frontier rather than being a routine application.

27.4 Priority D: mixed archimedean–radical invariants

The normalized radical $E = w^\vartheta$ has exact degree d and binomial minimal polynomial, so the Kummer field $\mathbb{Q}(E, \zeta_d)$ is completely explicit. The questions worth pursuing are whether the logarithmic relation $\log E \log 2 - \log w \log 3 = 0$ forces an impossible regulator or height relation there; whether simultaneous normalization at both outputs constrains the two normalized degrees through field intersections or discriminants beyond the known gcd condition; whether all complex branches $\log(\zeta_d^j E) = \log E + 2\pi i j/d$ can populate a multi-logarithm determinant with algebraic

entries of controlled height; and whether a p -adic logarithm branch is compatible with the real identity after raising to the normalized degree. A warning is essential, and it is the same one that closed Section 26.5: merely taking norms or conjugates of $X^d - c/3^a$ reproduces known relations, so the new invariant must mix conjugation with the special logarithmic direction.

27.5 Priority E: discharge and sharpen the determinant hypotheses

The algebraic and combinatorial determinant hierarchy has now been formalized and accepted:

ArbitraryNodeDeterminant, ExponentialPolynomialZeros,
OrderedChamberPositivity, MixedExponentSemigroup,
SemigroupDeterminant, FallingFactorialBound,
SemigroupWeightBound, ArbitraryNodeRepulsion, SemigroupGapBound.

They are all on the aggregate import surface. [\[kernel-verified Lean\]](#) The two quantitative endpoints retain their analytic inputs as explicit hypotheses, `RpowDividedDifferenceMeanValue` and `NewtonFactorization`; those propositions are not asserted by the corpus. The priority is therefore no longer to translate the algebra, but to prove these named hypotheses in the required uniform form, or to replace them by a stronger factorization whose archimedean loss is small enough to interact with only $O(\log X)$ nodes. The exact tropical initial-form criterion of `TropicalInitialForm` is a useful prerequisite on the nonarchimedean side, but it does not supply the cancellation lower bound that this program needs.

27.6 Priority F: force a third algebraic exponential

The denominator-controlled theorem in `RationalThirdBase` is a precise gateway: for an auxiliary integer q containing a prime outside $\{2, 3\}$ it suffices, in several forms, to prove that q^x is rational with denominator supported by the known outputs, after which monomial injectivity and the three-base determinant force $x \in \mathbb{Q}$ and hence $x \in \mathbb{Z}$. What this would buy is a complete proof. The obstacle is the one identified in Section 26.7: straight monomials cannot introduce a new prime direction, so candidates must come from norms, resultants, or algebraic combinations arising in the primitive radical field, whose x th powers would then have to simplify.

27.7 Priority G: five and sharp six exponentials

Formalizing the five exponentials theorem, or Roy’s sharp six exponentials theorem [8], would add genuinely new restrictions — for instance transcendence of $e^{\gamma x}$ for algebraic $\gamma \neq 0$ — and would extend the rational-third-base result from a controlled denominator to an unrestricted rational output. The effort is comparable to the completed Gelfond–Schneider port, which is encouraging; the risk is that it be undertaken without a specific downstream lemma in mind, since as Section 26.7 shows these theorems constrain b^x and not the divisibility facts the conjecture actually needs.

27.8 Priority H: larger kernel-replayed regions and effective measures

The zero-free region $2^x < 4096$ is kernel-verified, while the current software frontier is 2^{27} ([\[software computation\]](#)) in Section 40. A source draft aiming at 2^{14} has verified arithmetic rows but no

accepted assembled theorem, so the next target is a generated, kernel-replayed certificate beyond 4096, with 16384 a useful engineering milestone rather than an established bound. Every descent constant improves by bookkeeping once the region grows, so the payoff is broad but shallow. The bottleneck is a compact kernel-replayable proof object: adaptive convergent enclosures, chunked certificate modules, and tiers shared along ranges. In the same vein, effective irrationality measures for ϑ from Páde or hypergeometric methods would replace per-value certificate tiers with a uniform polynomial criterion and likely extend verified regions by orders of magnitude. Their limitation is sharp: they say nothing directly about the generator $\beta = \log_2 M$ for unknown M , which would require an effective measure derived from the log-product relation (99) and therefore lies beyond ordinary linear forms in logarithms.

27.9 Priority I: computational reconnaissance with exact certificates

Finite scanning is evidence, not a path to infinity, but computation remains valuable in three focused roles: searching candidate determinant shapes for large systematic 2- or 3-adic valuations, feeding Priority A and Priority B; enumerating small primitive radical patterns (w, d, a, c) and recording local obstructions with exact rational inequalities rather than floating-point rejection; and rejecting candidates early using the primitive-pair conditions — no matched depth, no common perfect power, no matched affine decomposition — while recording near misses keyed by odd core, denominator, and rational-power index. The goal must be pattern discovery and theorem generation; the difficulty is the discipline not to spend the budget on extending a decimal barrier that proves nothing new.

27.10 Priority J: the colossally abundant interface

Formalizing the original application — that the rational-power form of the conjecture implies that quotients of consecutive colossally abundant numbers are prime [1] — would connect the corpus to its historical source and give the whole development a visible consequence. It is logically downstream of everything above, and the work is combinatorial and definitional rather than transcendental, so it should follow the core structure modules rather than compete with them.

28 Lean formalization blueprint

The new arguments of Sections 19, 20, 22, 23 and 12 were written so that each analytic step is a reusable lemma rather than a step inside one monolithic proof. The intended module layout sits under `ExponentialIdentities/TwoBaseIntegerExponent/` alongside the existing kernel-verified corpus of Section 4.

Since the first draft of this blueprint was written, all of the modules described below have been accepted by the Lean kernel and placed on the aggregate import surface `ExponentialIdentities.lean`. The formal determinant chain deliberately exposes its two remaining analytic inputs — `RpowDividedDifferenceMeanValue` and `NewtonFactorization` — as hypotheses of the quantitative endpoint theorems. Thus the algebraic identities and every deduction from those hypotheses are kernel-verified, but the corpus does not assert the hypotheses themselves. The descriptions below record the implemented architecture rather than a pending plan.

28.1 ArbitraryNodeDeterminant

Status: [\[kernel-verified Lean\]](#) — `ArbitraryNodeDeterminant.lean`, with the analytic mean-value input exposed in the later module `ArbitraryNodeRepulsion`.

This module formalizes the ordinary-power determinant and the arbitrary-node repulsion theorem of Section 19. Its dependencies are:

1. a determinant/divided-difference identity for the matrix whose columns are $1, x, \dots, x^{r-1}, f(x)$ evaluated at the nodes, so that the determinant factors as the Vandermonde product times the r th divided difference of f ;
2. a generalized mean value theorem for divided differences at distinct ordered nodes, which converts that divided difference into a derivative value at an interior point;
3. the existing derivative tower `rpowDerivativeFamily` and the falling-power coefficient `fallingRpowCoeff`, both already available in `HigherDifference`;
4. the integer-entry determinant lower bound: a matrix all of whose entries are integers has integer determinant, so a nonzero determinant has absolute value at least one.

The accepted module splits exactly along those four lines. It defines `nodeMatrix` and its integer shadow `nodeIntMatrix`, proves that candidate nodes make every entry an integer, derives `one_le_abs_det_nodeMatrix` from that, and then isolates the algebra of the final step in a hypothesis-carrying form: the analytic factorization $\det = V \cdot dd$ and the divided-difference bound $|dd| \leq B$ are inputs, and the conclusion is the repulsion inequality. A target statement for the fully analytic version is:

```
theorem vandermonde_lower_bound_of_twoBaseNaturalCandidates
  {r : ℕ} (hr : 2 ≤ r)
  (m : Fin (r + 1) → ℕ)
  (hstrict : StrictMono m)
  (hcand : forall i, TwoBaseNaturalCandidate (m i)) :
  (r.factorial : ℝ) / |fallingRpowCoeff logThreeDivLogTwo r| *
    (m 0 : ℝ) ^ ((r : ℝ) - logThreeDivLogTwo)
  <= vandermondeProduct r m
```

Listing 2: Schematic fully analytic arbitrary-node signature.

The intermediate form accepted by the kernel is the hypothesis-carrying one, which is worth stating separately because it is reusable for any f whose divided differences can be bounded:

```
theorem inv_le_vandermondeProduct_of_dividedDifference (r : ℕ)
  (m : Fin (r + 1) → ℕ) (hstrict : StrictMono m)
  (hcand : forall i, TwoBaseNaturalCandidate (m i)) {dd B : ℝ}
  (hB : 0 < B)
  (hdet : (nodeMatrix r m).det = vandermondeProduct r m * dd)
  (hdd_bound : |dd| ≤ B) (hne : (nodeMatrix r m).det <> 0) :
  1 / B ≤ vandermondeProduct r m
```

Listing 3: Accepted packaged repulsion step.

Working with a `Fin`-indexed strictly monotone family and an explicit pair product is easier than working with a `Finset` of nodes; the `Finset` form is recovered at the end through the monotone enumeration `Finset.orderEmbOfFin`.

28.2 GeneralizedVandermonde

Status: [\[kernel-verified Lean\]](#) — the zero-count/nonvanishing engine is accepted in `ExponentialPolynomialZeros`, and ordered-chamber positivity is accepted in `OrderedChamberPositivity`. Appendix D gives the paper proof and its Wronskian interpretation.

The positivity theorem for generalized Vandermonde determinants with arbitrary real exponents is useful well beyond this conjecture, and two formalization routes were available.

Wronskian route. Prove the Wronskian determinant formula of Appendix D directly, then invoke or develop the criterion that positive initial Wronskians produce an extended complete Chebyshev system and hence positive ordered evaluation determinants. This needs a small Chebyshev-system API that `mathlib` does not yet have.

Zero-count route. After the substitution $x = e^t$, prove by induction that a nonzero exponential polynomial with $r + 1$ distinct real exponents has at most r real zeros. Deduce nonvanishing of the evaluation determinant, and if the sign is wanted, obtain it by continuity and a coalescing-node limit.

The decisive observation for the present application is that only *nonvanishing* is needed downstream. The determinant in question has integer entries, and a nonzero integer already has absolute value at least one; the sign never enters the repulsion inequality. That removes the coalescing-limit argument from the critical path. The formalization first took the zero-count route in `ExponentialPolynomialZeros`; the independent accepted module `OrderedChamberPositivity` now supplies the strict sign as well.

28.3 ExponentialPolynomialZeros

Status: [\[kernel-verified Lean\]](#) — `ExponentialPolynomialZeros.lean`.

This module carries the zero-count induction. Write $g(t) = \sum_{j \leq n} c_j e^{\lambda_j t}$ with $\lambda_0 < \dots < \lambda_n$. Normalizing by $e^{-\lambda_0 t}$ preserves the zero set; differentiating kills the constant term and lowers the number of exponents by one; Rolle’s theorem between consecutive zeros supplies the zeros of the derivative. Induction on n then forces every coefficient to vanish.

Two Lean engineering difficulties dominate the effort, and neither is visible in the paper proof.

First, the Rolle step must produce n *distinct and ordered* interior points from $n + 1$ ordered zeros. `Mathlib’s exists_hasDerivAt_eq_zero` yields one interior point per gap, but nothing about how the points from different gaps compare. The implementation chooses them with `choose`, then proves

strict monotonicity by chaining $y_i < x_{i+1} \leq x_j < y_j$ through the monotonicity of the node family. When the hypothesis arrives as a `Finset` of at least $n + 1$ reals rather than an indexed family, the monotone enumeration `Finset.orderEmbOfFin` converts it, after cutting the set down to exactly $n + 1$ elements with `Finset.exists_subset_card_eq`.

Second, the induction reindexes. The derivative of an $(m + 1)$ -term exponential polynomial is an m -term one whose exponents are $\lambda_{j+1} - \lambda_0$ and whose coefficients are $c_{j+1}(\lambda_{j+1} - \lambda_0)$, so every statement has to be transported along `Fin.succ`, and the induction has to be set up with the order quantified first so that `induction` may generalize over exponents, coefficients and nodes simultaneously. In the accepted module this is done through a private auxiliary theorem of the shape `forall (n : N) (lam c : Fin (n + 1) -> R), ...`, from which the user-facing statements are one-line consequences.

```
theorem card_le_of_expPoly_eq_zero {n : N} (lam c : Fin (n + 1) -> R)
  (hlam : StrictMono lam) (hc : exists j, c j <> 0) (S : Finset R)
  (hzero : forall t, t in S -> expPoly lam c t = 0) :
  S.card <= n

theorem det_rpowVandermonde_ne_zero {n : N} (x lam : Fin (n + 1) -> R)
  (hx : StrictMono x) (hx0 : 0 < x 0) (hlam : StrictMono lam) :
  (rpowVandermonde x lam).det <> 0
```

Listing 4: Zero-count theorem and its determinant form.

A vanishing determinant yields a nonzero kernel vector, via `Matrix.exists_mulVec_eq_zero_iff`, and it is a real combination of the functions $x \mapsto x^{\lambda_j}$ vanishing at $n + 1$ distinct positive nodes. So the determinant statement follows from the zero-count statement.

28.4 MixedExponentSemigroup and SemigroupDeterminant

Status: [\[kernel-verified Lean\]](#) — `MixedExponentSemigroup.lean` and `SemigroupDeterminant.lean`.

The pair-indexed exponent is the whole of the definitional layer:

```
noncomputable def mixedExponent (ab : N x N) : R :=
  (ab.1 : R) + (ab.2 : R) * logThreeDivLogTwo
```

and the module proves, in this order:

- injectivity of `mixedExponent`, from irrationality of ϑ ;
- integrality of $m^{a+b\vartheta} = m^a (m^\vartheta)^b$ for every candidate m , packaged as a natural-number witness `candidateRpowNat`;
- closure under addition, so that the image is an `AddSubmonoid R`;
- the elementary comparisons $a + b \leq a + b\vartheta \leq a + 2b$ and $a + b\vartheta \leq 3q$ for $a, b \leq q$, which are what the counting arguments actually consume.

`SemigroupDeterminant` then builds the matrix $m_i^{\lambda_j}$ for candidate nodes and mixed exponents, shows every entry is a natural number, and concludes integrality and the lower bound $|\det| \geq 1$ whenever the determinant is nonzero — the nonvanishing being imported from `ExponentialPolynomialZeros`.

```
theorem one_le_abs_det_semigroupMatrix (r : N) (m : Fin (r + 1) -> N)
  (hcand : forall i, TwoBaseNaturalCandidate (m i))
  (ab : Fin (r + 1) -> N x N)
  (hne : (semigroupMatrix r m ab).det <> 0) :
  1 <= |(semigroupMatrix r m ab).det|
```

Listing 5: Semigroup determinant integrality.

It is deliberately unnecessary to define the globally sorted semigroup here. The determinant estimate accepts an arbitrary finite family of pairs together with a proof that their mixed exponents are strictly increasing; the sorted enumeration is a convenience of the paper proof, not a requirement of the formal one.

28.5 FallingFactorialBound

Status: [\[kernel-verified Lean\]](#) — `FallingFactorialBound.lean`, on the aggregate import surface.

This module carries the two elementary estimates behind every Newton-row bound of Section 20: the falling-coefficient bound $|(\lambda)_i|/i! \leq 2^{\lambda+i}$ for $\lambda \geq 0$, and the strip bound $x^{\lambda-i} \leq 2^\lambda N^{\lambda-i}$ for $N \leq x \leq 2N$, $N \geq 1$. The proof of the first compares each factor $|\lambda - k|$ with $m + 1 + k$ where $m = \lfloor \lambda \rfloor$, identifies $\prod_{k < i} (m + 1 + k)$ with $i!(m+i)$, and bounds the binomial coefficient by the row sum 2^{m+i} . Their product is the single inequality consumed downstream:

```
theorem abs_fallingRpowCoeff_div_factorial_mul_rpow_le
  {N x lam : R} (hN : 1 <= N) (h1 : N <= x) (h2 : x <= 2 * N)
  (hlam : 0 <= lam) (i : N) :
  |fallingRpowCoeff lam i| / (i.factorial : R) * x ^ (lam - (i : R))
  <= (2 : R) ^ (2 * lam + (i : R)) * N ^ (lam - (i : R))
```

Listing 6: Newton-row entry bound.

The determinant estimate then follows from `Matrix.det_apply` or, better, from a generic Leibniz-bound lemma worth adding once:

if every entry in row i , column j is bounded by a separable quantity $A_i B_j$, then the determinant is at most $(r + 1)! \prod_i A_i \prod_j B_j$.

28.6 SemigroupWeightBound

Status: [\[kernel-verified Lean\]](#) — `SemigroupWeightBound.lean`.

For the explicit logarithmic-square corollary the full asymptotic count of the ordered semigroup is optional. What is needed is the square-box construction of Section 20 in *existence form*, which sidesteps real-valued order statistics entirely:

- the $(q+1)^2$ pairs in $[0, q]^2$ have pairwise distinct mixed exponents;
- every such exponent is at most $3q$, using $1 < \vartheta < 2$;
- taking $q = \lfloor \sqrt{r} \rfloor$ makes $r+1$ pairs fit, with total weight at most $6(r+1)\sqrt{r}$;
- dividing by $R_r = r(r+1)/2$ gives the deficit ratio $\Sigma_r/R_r \leq 12/\sqrt{r}$.

The module realizes the injective family explicitly by base- $(q+1)$ digits, $j \mapsto (j \bmod (q+1), j \div (q+1))$, so no sorting is performed and no canonical ordered semigroup is ever constructed. The statements are:

```
theorem exists_mixedExponent_family (r : N) :
  exists ab : Fin (r + 1) -> N x N, Function.Injective ab and
    (forall j, mixedExponent (ab j) <= 6 * Real.sqrt (r : R)) and
    (sum j, mixedExponent (ab j)) <= 6 * ((r : R) + 1) * Real.sqrt (r : R)
```

Listing 7: Square-box weight bound, existence form.

The module also contains the abstract counting bridge that turns any span lower bound into a cardinality bound, and which serves the ordinary-power corollary of Section 19 equally well: if $S \subseteq [N, N+H]$ and every $r+1$ elements of S contain two points at distance at least $L > 0$, then $\#S \leq r + rH/L$. The proof partitions $[N, N+H]$ into $\lfloor H/L \rfloor + 1$ half-open blocks of length L , observes that each block meets S in at most r points, and applies pigeonhole.

28.7 SemigroupDyadicBound

Status: [\[kernel-verified Lean\]](#) — implemented in `SemigroupGapBound`; the final cardinality theorem is conditional on the explicit `NewtonFactorization` hypothesis.

This layer combines the span theorem with the consecutive-window summation and produces the explicit $O((\log X)^2)$ dyadic bound. The accepted endpoint is `card_le_ten_thousand_max_log_sq`; schematically it states:

```
theorem card_twoBaseNaturalCandidates_Icc_le_log_sq (X : N) (hX : 1 <= X) :
  ((Finset.Icc X (2 * X)).filter TwoBaseNaturalCandidate).card
  <= Nat.ceil (10000 * max 1 (Real.log X) ^ 2)
```

Listing 8: Schematic explicit dyadic bound.

A cleaner all-natural statement may replace the real logarithm by the binary ceiling logarithm, obtaining a larger but much easier constant times `Nat.clog 2 X ^ 2`. Since the repository already has logarithmic-square counting infrastructure in `MultiplicativeRank`, its finite-set definitions and cardinality lemmas can be reused while keeping the proof mechanism entirely independent — which is the point of proving the same bound twice.

28.8 BlockGrowthReformulations

Status: [\[kernel-verified Lean\]](#) — `BlockGrowthReformulations.lean`.

The exact conditional block formula proved in the corpus makes the dyadic block count $B(T) = \#\{\text{candidates in } [2^T, 2^{T+1})\}$ either identically 1 or bounded below by a positive multiple of T . That dichotomy collapses a whole family of apparently weaker growth hypotheses onto the conjecture itself, and the module states all four equivalences together so that a future argument may pick whichever is most convenient:

```
theorem alaogluErdosConjecture_block_growth_reformulations :
  (AlaogluErdosConjecture <=> exists C : N, forall T : N, dyadicBlockCount T <= C) and
  (AlaogluErdosConjecture <=>
    (fun T : N => (dyadicBlockCount T : R)) =o[atTop] fun T : N => (T : R)) and
  (AlaogluErdosConjecture <=>
    forall eps : R, 0 < eps ->
      frequently T in atTop, (dyadicBlockCount T : R) < eps * (T : R)) and
  (AlaogluErdosConjecture <=> forall T : N, dyadicBlockCount T = 1)
```

Listing 9: Block-growth reformulations.

The practical value is negative and worth stating plainly: proving mere boundedness, or even vanishing lower density along a subsequence, is exactly as hard as the conjecture. No softening of the counting target buys anything.

28.9 CompactOrbit

Status: [\[kernel-verified Lean\]](#) — `CompactOrbit.lean`.

This module states the conjecture as the nonexistence of a synchronized integral orbit on the compact arc $\{(2^t, 3^t) : 0 \leq t < 1\}$. The orbit point is

$$(M^k / 2^{\lfloor k\beta \rfloor}, A^k / 3^{\lfloor k\beta \rfloor}),$$

both coordinates divided by a power of their own base with the *same* integer exponent $\lfloor k\beta \rfloor$; that shared exponent is what makes the reformulation nontrivial.

```
theorem alaogluErdosConjecture_iff_no_compactOrbit :
  AlaogluErdosConjecture <=>
    forall M A : N, forall beta : R,
      1 < M -> 1 < A -> Irrational beta -> 0 < beta ->
        exists k : N, 1 <= k and
          synchronizedOrbitPoint M A beta k not_in compactExponentialArc
```

Listing 10: Compact-orbit reformulation.

The forward direction routes through the canonical primitive generator, so the produced β is the least noninteger solution and M, A are its normalized outputs. The converse needs neither $M, A > 1$ nor $\beta > 0$: arc membership by itself already pins $M = 2^\beta$ and $A = 3^\beta$. Formalizing that asymmetry was the only delicate part.

28.10 LogProductIndependence

Status: [\[kernel-verified Lean\]](#) — `LogProductIndependence.lean`. The independence property is an explicit hypothesis, never asserted; the implication from it is accepted.

This module records the prime-log-product criterion of Section 12. The open hypothesis is never asserted; it is a named `Prop` that appears only as an explicit argument, so the exported theorem is unconditional Lean about a conditional mathematical statement.

```
def PrimeLogProductIndependence : Prop :=
  forall (S : Finset N) (c : N -> N -> Z),
    (forall p, p in S -> Nat.Prime p) ->
      (sum p in S, sum q in S, (c p q : R) * (Real.log p * Real.log q)) = 0 ->
        forall p, p in S -> forall q, q in S -> c p q + c q p = 0

theorem alaogluErdosConjecture_of_primeLogProductIndependence
  (hind : PrimeLogProductIndependence) : AlaogluErdosConjecture
```

Listing 11: Prime-log-product criterion.

The supporting lemma is the prime-factorization expansion $\log n = \sum_{p|n} v_p(n) \log p$, after which the relation $M^{\log 3} = A^{\log 2}$ becomes a vanishing integer combination of the products $\log p \log q$; symmetrized coefficientwise vanishing then forces M to be a power of 2 and the exponent to be an integer.

28.11 Proof-audit discipline

The discipline that governs every status label in this report is mechanical, and it is worth restating because several modules above are new. Each headline theorem must enter the aggregate import surface `ExponentialIdentities.lean` and be checked with `#print axioms`, so that the trust boundary of each result is visible rather than inferred. A result stays **[Lean draft, not yet accepted]** or **[paper proof in this report]** until a full build succeeds; the existence of a `sorry`-free source file is not verification, and neither is a successful build of an earlier revision. Numerical constants appearing in the tables of this report are explanatory only — the formal statements use exact expressions or safe rational constants such as $1/24$ and 10000 , chosen so that the arithmetic side conditions close by `norm_num` rather than by a tuned estimate.

29 Comparison of methods

29.1 Prime-valuation rank versus interpolation geometry

The module `MultiplicativeRank` proves that any three candidates have linearly dependent prime-exponent vectors, of rank at most two. It does so by a three-base, six-exponentials determinant: if three candidate bases were multiplicatively independent, integrality of all three ϑ th powers would force ϑ rational, contradicting its irrationality. A finite candidate set therefore lies in a rational plane inside the free vector space on primes. Choosing two prime coordinates injectively encodes the set and yields an explicit logarithmic-square count.

The semigroup determinant of Section 20 proves a logarithmic-square dyadic bound by an entirely different mechanism.

Prime-valuation method	Semigroup determinant method
Uses prime factorization of candidate bases	Uses only integer values on $y = x^\vartheta$
Global multiplicative rank at most two	Local total positivity and interpolation
Invokes a six-exponentials consequence	Uses elementary determinant integrality once ϑ is fixed
Produces coordinates in two prime directions	Produces repulsion for arbitrary ordered nodes
Naturally global in magnitude	Naturally local in additive intervals
Kernel-verified [kernel-verified Lean]	Paper theorem [paper proof in this report] ; hypothesis-carrying Lean endpoint [kernel-verified Lean]

Their agreement at logarithmic-square scale is suggestive rather than coincidental. Both methods see a two-dimensional structure, but they see different two-dimensional structures: prime-exponent rank two on one side, mixed-exponent lattice dimension two on the other. Neither refinement alone has closed the gap, and a future proof may well need to combine the two dual lattices rather than push either one further in isolation.

29.2 Exact conditional rank one over two generators

Conditional on failure of the conjecture, the candidate monoid is not merely of rank at most two — it is exactly free on 2 and M . In prime-valuation space its points form the nonnegative lattice generated by the prime-exponent vectors of 2 and M . In mixed-exponent evaluation space, the values factor across a two-dimensional grid indexed by the pair (n, k) with value $2^n M^k$. This exact grid is far stronger than the unconditional rank theorem, and yet every presently known consequence remains compatible with it. That compatibility is the honest measure of how much room is left.

The strongest likely synthesis is a determinant whose rows use primitive-generator coordinates and whose columns use mixed exponents. Such a matrix would be simultaneously:

- integral, hence subject to the $|\det| \geq 1$ lower bound;
- totally positive after ordering by magnitude, hence nonvanishing for free;
- explicitly factored into powers of 2, 3, M , A ;
- constrained by affine primitivity and by simultaneous output normalization.

No current module combines all four properties. `MultiplicativeRank` has the first and third; the semigroup modules have the first and second; `PrimitiveGenerator` and `OutputNormalization` have the fourth. Building the module that has all four is the most concrete open formalization task suggested by this report.

29.3 Finite differences as confluent determinants

Forward differences are the equally spaced specialization of interpolation, and Wronskians are its confluent limit as nodes coalesce. The determinant hierarchy of Sections 19 and 20 therefore provides

a single conceptual envelope for four things that the corpus currently treats separately:

- second-, third-, and arbitrary-order finite differences;
- arbitrary-node ordinary-power determinants;
- mixed-monomial generalized Vandermonde determinants;
- confluent determinants using derivatives or repeated coordinate directions.

The corresponding formalization strategy is to build a general interpolation-determinant layer once and recover **HigherDifference** as a specialization of it, rather than to keep the two developments parallel. That would remove duplicated real analysis and, more usefully, make experimental determinant choices cheap to certify: a new choice of nodes and exponents would need only a divided-difference bound, with integrality and nonvanishing supplied by the shared layer.

30 Conclusions

The Alaoglu–Erdős conjecture (Conjecture 2.1) remains open. Nothing in this report proves it, and no combination of the results collected here proves it. What has changed is that the reason can now be stated exactly rather than impressionistically.

The kernel-verified core is broad and its themes are complementary. Interpolation determinants settle three independent bases; Gelfond–Schneider excludes algebraic irrational exponents; exact certificates create a growing zero-free region, currently $2^x < 4096$; unique factorization yields a canonical primitive odd core and, conditionally, an exact solution monoid; arbitrary-order forward differences force sharp local additive sparsity, with the rational three-term criterion $h^{400} \leq n^{83}$ as its effective form; leastness gives exact affine primitivity; and unconditional equivalences with fractional-part gaps and localized radicals identify the entire remaining distance to the conjecture.

Conditional on a single counterexample, the corpus determines the whole picture. A counterexample would not produce a messy exceptional set that density or semigroup arguments might squeeze; it would produce the perfectly rigid free monoid

$$\mathcal{S} = \mathbb{N}_0 + \mathbb{N}_0\beta, \quad \mathcal{C} = \{2^n M^k : n, k \in \mathbb{N}_0\},$$

whose block counts, cumulative asymptotics, fractional parts, rotation statistics, rational-power index, radical degree, and simultaneous output normalization are all explicit and internally consistent with every verified gap and density bound. The surviving obstruction is correspondingly concentrated:

$$\text{exclude } w^{\log_2 3} \in \mathbb{Z}[1/3] \text{ for every primitive odd } w > 1.$$

That single exclusion is equivalent to, and exactly as open as, the original conjecture.

The new determinant hierarchy sharpens the picture quantitatively. The arbitrary-node theorem of Section 19 shows that candidate lattice points cannot cluster freely: every Vandermonde product over candidate nodes obeys a power-law lower bound. The mixed-exponent semigroup of Section 20 upgrades this to arbitrarily high-order positive integer determinants with a repulsion exponent tending to one, and yields a second, conceptually independent logarithmic-square counting theorem.

The most valuable outcome, however, is the quantified failure point of Section 22. Pure archimedean interpolation needs about $(\log X)^2$ nodes to operate at scale X , whereas the exact conditional

monoid supplies only about $\log X$ of them. The method misses by exactly one factor of $\log X$ — not by a constant, and not by an epsilon. This makes future work sharper by exclusion: improving constants, tightening the falling-factorial bound, or extending finite searches cannot close a gap of that shape. The plausible routes are precisely those capable of supplying one additional logarithm of determinant strength: 2- or 3-adic divisibility used simultaneously with the archimedean estimate, a synchronized-denominator counting theorem on the compact exponential arc of Section 23, or a mechanism forcing a third algebraically independent exponential direction.

Section 12 isolates a different sufficient condition, and it is worth separating from the rest because it does not live at the four-exponentials boundary at all. If the products $\log p \log q$ of prime logarithms are linearly independent over $\mathbb{Q}$, the conjecture follows by an argument that is elementary once the hypothesis is granted. That independence is a consequence of Schanuel’s conjecture and is not known unconditionally — indeed even the irrationality of $(\log 3)^2/(\log 2)^2$ is open — but it is a different classical conjecture, with a different literature and different partial results. Having two inequivalent sufficient conditions, one transcendence-theoretic and one arithmetic, is more informative than having one.

The near-term program is the one consolidated in Section 27: finish the exact conditional consequences, in particular the functional-analytic extension of blockwise equidistribution to arbitrary continuous tests and interval indicators; exploit the uniform higher-difference theorem combinatorially, allowing the order to grow with height; kernel-replay the 2^{20} region as a compact certificate; use the explicit radical field for a genuinely number-field argument; and pursue auxiliary-function methods tailored to integrality, where all corner values are positive integers and the primitive pair is canonical. Alongside those, the immediate formalization work is the module list of Section 28: bring the nine existing drafts through a full kernel build and onto the aggregate import surface, then attempt the combined determinant identified in Section 29.

A complete proof still requires a new idea. But the target is now narrow, the conditional structure is exact, the failure mode of the best available elementary method is measured rather than guessed, and the surrounding formal infrastructure is strong enough to test the next idea precisely.

31 Continuation: executive answer, scope, and trust boundaries

The sections that follow were produced as a continuation of the preceding material and are merged here unchanged in substance. They supply a symbolic-dynamical model of a hypothetical counterexample, an effective replacement for the elementary continued-fraction estimate, exact all-support valuations of the block determinants, and the sharp constants that govern how much of the archimedean budget term-wise divisibility can supply. Section 21 states the unconditional ceiling that underlies all of them and records the two constants this part corrects.

31.1 The requested goal and the outcome

The requested primary goal was to continue the attached research program and prove the following statement.

Conjecture 31.1 (Alaoglu–Erdős). *If $x \in \mathbb{R}$ and $2^x, 3^x \in \mathbb{Z}$, then $x \in \mathbb{Z}$.*

I did not find an unconditional proof. In particular, none of the arguments below converts

$$\log(2^x) \log 3 - \log(3^x) \log 2 = 0$$

into a forbidden linear form in logarithms of algebraic numbers. The equality is bilinear in logarithms, and normalizing it gives exactly the unresolved 2×2 determinant behind the four exponentials conjecture. Baker-type estimates control every *nonzero* linear form that occurs after fixing integer coefficients; they do not prove that this special bilinear form cannot vanish.

The continuation nevertheless changes the technical map in useful ways. It obtains an exact symbolic-dynamical formulation, replaces an elementary continued-fraction estimate by an effective polynomial one, computes the full leading p -adic content of ordinary block Vandermondes, and gives sharp continuum formulas for the termwise divisibility of two much larger determinant constructions. These calculations do more than say that a method “seems too weak”: they quantify the leading constant that is available and isolate the precise kind of extra cancellation a successful proof would have to produce.

31.2 The principal new results

Result	Content	Status
Synchronized Sturmian orbit	A counterexample is equivalent to one irrational mechanical word simultaneously controlling the dyadic and triadic denominator jumps of two compact rational recurrences.	[paper proof in this report]
Effective irrationality of β	Baker–Wüstholz/Matveev gives $ \beta - p/q \geq c(M)q^{-\mu(M)}$ and hence polynomial, rather than exponential, recurrence bounds for continued-fraction denominators.	[classical/open back-ground] plus [paper proof in this report]

Result	Content	Status
Ordinary Vandermonde audit	The valuations at every prime in the fixed supports of M and A are computed to leading order. The product construction has universal ceiling $1 - \log 2 / (3 \log 6) = 0.871049 \dots$	[paper proof in this report]
Gini assignment formula	For the balanced mixed-semigroup determinant, the exact termwise all-prime divisor fraction is $1 - \frac{3}{8} \sum_p \mathfrak{G}(c_p) + o(1)$ for an explicit norm $\mathfrak{G}$.	[paper proof in this report]
Universal 4/5 ceiling	The primitive condition forces $\sum_p \mathfrak{G}(c_p) \geq 8/15$. Thus entrywise/min-plus divisibility can supply at most 80% of the leading block determinant height.	[paper proof in this report]
Cross-coprime constants	If $\gcd(MA, 6) = 1$, the all-prime ceiling is 7/10, and the places 2 and 3 alone contribute exactly 3/10 in the continuum limit.	[paper proof in this report]
Full-grid transport ceiling	On the complete triangular (n, k) grid, any termwise all-prime divisor is bounded by the countermonotone energy $\pi/8$, while the largest determinant term has energy $1/2$; the ratio is $\pi/4$.	[paper proof in this report]
HCIZ reformulation	The full-grid archimedean problem is expressed by the Harish–Chandra–Itzykson–Zuber integral, identifying a concrete free-energy inequality as the remaining analytic target.	[classical/open background] plus [paper proof in this report]

31.3 Evidence labels used in this part

This part uses the report-wide convention of Section 3.1, with two additions. [repository build reported] marks a statement for which a repository commit or build log reports success but which was not independently replayed by the author of the continuation; it is repository metadata, not a kernel audit. [proposed target] marks a statement proposed as a precise target rather than asserted. Where the continuation wrote [classical/open background] for standard literature the report-wide label is used unchanged.

31.4 What the determinant ceilings do and do not prove

A ceiling such as 4/5 has a specific meaning. Given an integer determinant $D = \det(E_{ij})$, define at each prime p the minimum valuation of a Leibniz term,

$$\tau_p = \min_{\sigma} \sum_i v_p(E_{i, \sigma(i)}).$$

Every term, and therefore D , is divisible by p^{τ_p} . The product $\prod_p p^{\tau_p}$ is the *termwise* or *tropical* divisor. It ignores extra valuation caused by cancellation among several terms attaining the same

minimum. The 4/5 theorem says that this automatic divisor cannot occupy more than four-fifths of the leading archimedean height for the balanced block construction. It does *not* say that the true valuation $v_p(D)$ is at most the tropical minimum; the opposite inequality holds, and extra cancellation may be large. Thus the result retires only the naïve entrywise-divisibility version of the method. It leaves a sharply defined cancellation problem.

32 Baseline inherited from the unified report

32.1 Notation

Set

$$\vartheta = \log_2 3.$$

The solution and candidate sets are

$$\mathcal{S} = \{x \geq 0 : 2^x, 3^x \in \mathbb{N}\}, \quad \mathcal{C} = \{m \in \mathbb{N} : m^\vartheta \in \mathbb{N}\}.$$

The map $x \mapsto 2^x$ is a bijection $\mathcal{S} \rightarrow \mathcal{C}$. Integer exponents give the obvious solutions $\mathbb{N}_0 \subseteq \mathcal{S}$ and candidates $2^{\mathbb{N}_0} \subseteq \mathcal{C}$.

Throughout the conditional sections, suppose the conjecture fails. The exact structure proved in the source report then supplies a least nonintegral solution $\beta > 0$ and primitive outputs

$$M = 2^\beta \in \mathbb{N}, \quad A = 3^\beta \in \mathbb{N}.$$

Every solution and candidate has a unique coordinate representation

$$\mathcal{S} = \{n + k\beta : n, k \in \mathbb{N}_0\}, \quad \mathcal{C} = \{2^n M^k : n, k \in \mathbb{N}_0\}. \quad (100)$$

The associated output is

$$(2^n M^k)^\vartheta = 3^n A^k.$$

We write

$$a = v_2(M), \quad b = v_3(A), \quad c = v_2(A), \quad d = v_3(M).$$

The inherited primitive-output theorem includes the crucial condition

$$\min(a, b) = 0. \quad (101)$$

It is important not to silently strengthen (101) to $a = b = 0$.

32.2 Inherited results used here

The following facts are taken from the attached unified report and its referenced Lean modules.

- [\[kernel-verified Lean\]](#) A nonintegral solution, if it exists, is transcendental. The constant ϑ is transcendental, by Gelfond–Schneider.
- [\[kernel-verified Lean\]](#) Failure gives the exact free rank-two monoid structure (100), including uniqueness of (n, k) .

- [\[kernel-verified Lean\]](#) In a dyadic exponent block $[T, T + 1)$, the solutions are exactly

$$T \quad \text{and} \quad T - \lfloor k\beta \rfloor + k\beta, \quad 1 \leq k \leq \left\lfloor \frac{T+1}{\beta} \right\rfloor.$$

Thus the block count is $1 + \lfloor (T+1)/\beta \rfloor$.

- [\[kernel-verified Lean\]](#) The primitive pair obeys (101); both outputs cannot simultaneously be $\{2, 3\}$ -smooth.
- [\[kernel-verified Lean\]](#) Mixed monomials $m^u(m^\vartheta)^v$ are integers for $(u, v) \in \mathbb{N}_0^2$, and generalized Vandermonde determinants built from distinct candidate nodes and distinct exponents $u + \vartheta v$ are nonzero integers.
- [\[kernel-verified Lean\]](#) The kernel-checked finite exclusion in the source report gives $x \geq 12$ for a nonintegral solution. The report separately records a directed-rounding software exclusion through $x \leq 26$.

The continuation does not reprove this large inherited corpus. It restates every inherited fact at the point where it enters a new argument.

32.3 Repository reconciliation as of August 21, 2026

The attached source is internally newer than one bibliographic pin printed near its end: its body and status matrix already discuss the enlarged algebraic-output and determinant module surface, although a repository item still names commit `5d3982ceb12644fdd3499569d9db5f62f9a31dc6`. During this continuation I checked the current public repository history. Commit `78f80c5d3f6f9afd64ecb72ddb154c85e2036bc1` reports a successful aggregate build with 67 imported modules and merges the later algebraic-output work. Subsequent commits add a literature compass and a short p -adic priority note; the latest snapshot consulted while writing was `ff9a9f5aa3a3f02599de93fc053e60918f269855`.

Status: [\[repository build reported\]](#)

The build claims in the preceding paragraph are taken from repository commit messages and files. I did not independently run all Lean jobs. The mathematical arguments newly stated below are therefore labelled [\[paper proof in this report\]](#), not [\[kernel-verified Lean\]](#).

32.4 The exact remaining logarithmic wall

Writing $M = 2^\beta$ and $A = 3^\beta$ gives

$$\beta = \frac{\log M}{\log 2} = \frac{\log A}{\log 3}$$

and hence

$$\log M \log 3 - \log A \log 2 = 0. \tag{102}$$

The four numbers $\log 2, \log 3, \log M, \log A$ are logarithms of algebraic numbers. Equation (102) is a vanishing 2×2 determinant. Four exponentials would exclude it because $2, 3$ are multiplicatively independent and $1, \beta$ are linearly independent over $\mathbb{Q}$, but four exponentials remains open. The new work below therefore pursues information that is invisible in the bare real logarithmic determinant: exact floor coding, effective separation, and valuations of large integer determinants.

33 A synchronized Sturmian model of a counterexample

33.1 Fractional-part normalization

Write

$$\beta = B + \alpha, \quad B = \lfloor \beta \rfloor \in \mathbb{N}_0, \quad 0 < \alpha < 1.$$

Since a nonintegral solution is transcendental, α is irrational. Define the rational localized multipliers

$$U = \frac{M}{2^B} = 2^\alpha \in \mathbb{Q} \cap (1, 2), \quad V = \frac{A}{3^B} = 3^\alpha \in \mathbb{Q} \cap (1, 3). \quad (103)$$

For $k \geq 0$, put

$$e_k = \lfloor k\alpha \rfloor, \quad s_k = e_{k+1} - e_k \in \{0, 1\}.$$

The binary sequence $s = s_0 s_1 s_2 \dots$ is the lower mechanical word of slope α . Define compact normalized orbits

$$R_k = \frac{U^k}{2^{e_k}} = 2^{\{k\alpha\}} \in [1, 2), \quad S_k = \frac{V^k}{3^{e_k}} = 3^{\{k\alpha\}} \in [1, 3). \quad (104)$$

Both orbits use the *same* fractional part and hence the same symbolic itinerary.

Theorem 33.1 (Synchronized dyadic–triadic Sturmian recurrence). *Assume a primitive nonintegral solution β exists. Then the rational sequences R_k, S_k in (104) satisfy*

$$R_0 = S_0 = 1, \quad R_{k+1} = \frac{UR_k}{2^{s_k}}, \quad S_{k+1} = \frac{VS_k}{3^{s_k}}, \quad (105)$$

where

$$s_k = 1 \iff \{k\alpha\} \geq 1 - \alpha \iff UR_k \geq 2 \iff VS_k \geq 3. \quad (106)$$

Moreover,

$$\sum_{j=0}^{N-1} s_j = \lfloor N\alpha \rfloor \quad (N \geq 1). \quad (107)$$

Proof. The recurrence follows by subtracting consecutive floor exponents:

$$R_{k+1} = \frac{U^{k+1}}{2^{e_{k+1}}} = \frac{UU^k}{2^{e_k + s_k}} = \frac{UR_k}{2^{s_k}},$$

and the same calculation with base 3 gives the formula for S_k . Since $e_{k+1} = e_k + 1$ exactly when $\{k\alpha\} + \alpha \geq 1$, the first equivalence in (106) holds. The identities $UR_k = 2^{\alpha + \{k\alpha\}}$ and $VS_k = 3^{\alpha + \{k\alpha\}}$ give the other two. Finally, (107) telescopes: $\sum_{j < N} s_j = e_N - e_0 = \lfloor N\alpha \rfloor$. $\square$

33.2 Sturmian consequences

The following facts are classical consequences of irrational mechanical coding [25, 26].

Corollary 33.2 (Aperiodicity, balance, and minimal complexity). *The word s in Theorem 33.1 is aperiodic. Any two factors of equal length contain numbers of 1's differing by at most one, and the number of distinct factors of length n is exactly $n + 1$.*

Aperiodicity has an elementary proof in this setting. If s were eventually periodic, then its limiting density of 1's would be rational. Equation (107) shows that this density is α , contradicting irrationality. Balance and factor complexity are the standard Sturmian characterization.

The point is not merely terminological. A hypothetical counterexample is no longer described only by two rational numbers U, V satisfying a logarithmic equality. It determines one minimal-complexity aperiodic word, and that same word must be realized by two rational piecewise-multiplicative systems with multiplicatively independent reset bases 2 and 3. This is a much more rigid synchronized object.

33.3 Exact reduced denominators

The primitive factor depths determine how much cancellation occurs in (104). Write

$$M = 2^a W, \quad W \text{ odd}, \quad A = 3^b Z, \quad 3 \nmid Z.$$

Because M is not a power of 2 and A is not a power of 3, one has $W, Z > 1$.

Proposition 33.3 (Denominator exponents and common jump word). *In lowest terms,*

$$R_k = \frac{W^k}{2^{q_k^{(2)}}}, \quad q_k^{(2)} = \lfloor k\beta \rfloor - ak = k(B - a) + \lfloor k\alpha \rfloor, \quad (108)$$

$$S_k = \frac{Z^k}{3^{q_k^{(3)}}}, \quad q_k^{(3)} = \lfloor k\beta \rfloor - bk = k(B - b) + \lfloor k\alpha \rfloor. \quad (109)$$

Consequently

$$q_{k+1}^{(2)} - q_k^{(2)} = B - a + s_k, \quad q_{k+1}^{(3)} - q_k^{(3)} = B - b + s_k.$$

At least one of the two sequences has baseline jump B , because $\min(a, b) = 0$.

Proof. Using $\lfloor k\beta \rfloor = kB + \lfloor k\alpha \rfloor$,

$$R_k = \frac{M^k}{2^{\lfloor k\beta \rfloor}} = \frac{2^{ak} W^k}{2^{\lfloor k\beta \rfloor}} = \frac{W^k}{2^{\lfloor k\beta \rfloor - ak}}.$$

The numerator is odd, so this is reduced. The base-3 calculation is identical. Taking successive differences gives the jump formulas, and the last claim is the inherited primitive condition (101). $\square$

Thus the same Sturmian bit s_k records whether *both* reduced denominators take their large or small permitted jump at time k . One denominator can have a shallower constant baseline, but the aperiodic fluctuation is synchronized exactly.

33.4 An exact equivalent exclusion principle

The source report already contains localization reformulations such as $2^\alpha \in \mathbb{Z}[1/2]$ and $3^\alpha \in \mathbb{Z}[1/3]$. The recurrence packages them into a symbolic statement.

Theorem 33.4 (Synchronized Sturmian exclusion is equivalent to Alaoglu–Erdős). *The Alaoglu–Erdős conjecture is equivalent to the nonexistence of data*

$$\alpha \in (0, 1) \setminus \mathbb{Q}, \quad U \in \mathbb{Q} \cap (1, 2), \quad V \in \mathbb{Q} \cap (1, 3)$$

such that, for the lower mechanical word $s_k = \lfloor (k+1)\alpha \rfloor - \lfloor k\alpha \rfloor$, the recurrences

$$R_0 = S_0 = 1, \quad R_{k+1} = UR_k/2^{s_k}, \quad S_{k+1} = VS_k/3^{s_k}$$

remain in $[1, 2) \cap \mathbb{Q}$ and $[1, 3) \cap \mathbb{Q}$, respectively, for every k , with $U = 2^\alpha$ and $V = 3^\alpha$. Equivalently, it suffices to rule out irrational α for which

$$2^\alpha \in \mathbb{Z}[1/2], \quad 3^\alpha \in \mathbb{Z}[1/3].$$

Proof. A counterexample $\beta = B + \alpha$ gives the data by Theorem 33.1. Conversely, if such α, U, V exist, choose an integer B at least as large as the powers of 2 and 3 occurring in the respective reduced denominators of U and V . Then

$$2^{B+\alpha} = 2^B U \in \mathbb{N}, \quad 3^{B+\alpha} = 3^B V \in \mathbb{N},$$

while $B + \alpha \notin \mathbb{Z}$. The recurrence identities are equivalent to $R_k = 2^{\{k\alpha\}}$ and $S_k = 3^{\{k\alpha\}}$, so they add an exact dynamic coding but no hidden hypothesis. $\square$

33.5 Cobham-type and Mahler-type diagnostics

Cobham’s theorem says that a sequence automatic in two multiplicatively independent integer bases is ultimately periodic [27, 28]. Sturmian words of irrational slope are not automatic in any integer base. Theorem 33.4 therefore suggests a tempting strategy: show that the dyadic rational recurrence makes s 2-automatic and that the triadic recurrence makes it 3-automatic.

That implication is *not* presently available. Rationality of the multiplier does not turn the compact rational state R_k or S_k into a finite-state object; the exact reduced denominators grow without bound. Cobham’s theorem becomes applicable only after proving a finite-kernel or finite-substitution property, and that is essentially a new arithmetic statement. The synchronized formulation is useful because it identifies this missing lemma cleanly rather than treating “automaticity” as a metaphor.

Research target 33.5 (Finite-kernel extraction). Show that simultaneous rational realizability of the same irrational mechanical word by (105) forces a finite 2-kernel and a finite 3-kernel, or derive a weaker pair of incompatible substitutive structures.

A Mahler-function variant is also conceivable: encode the jump word by $F(z) = \sum_{k \geq 0} s_k z^k$. Aperiodicity and the finite coefficient set already imply, by classical results of Fatou type, that F is transcendental over $\mathbb{Q}(z)$. A proof would have to derive incompatible functional equations from the two rational recurrences. No such equations were found here; the recurrence is multiplicative in the state, not self-similar in the index.

34 Effective logarithmic separation

34.1 From linear forms in two logarithms to an irrationality measure

The source report proves an elementary denominator-growth statement: a very good rational approximation to a hypothetical solution β would force a correspondingly large output, and consecutive continued-fraction denominators obey a single-exponential recurrence. Since β is itself a ratio of logarithms of fixed integers, Baker’s theory gives a substantially stronger conclusion.

Theorem 34.1 (Effective finite irrationality measure for the primitive generator). *Assume a primitive counterexample β exists, and put $M = 2^\beta$. There are effectively computable constants $c_M > 0$ and $\mu_M < \infty$ such that, for every $p \in \mathbb{Z}$ and $q \geq 2$,*

$$\left| \beta - \frac{p}{q} \right| \geq c_M q^{-\mu_M}. \quad (110)$$

One may choose μ_M bounded by an effectively computable absolute constant times $1 + \log M$, hence by an absolute constant times $1 + \beta$. The same conclusion follows from $\beta = \log A / \log 3$.

Proof. Consider the linear form

$$\Lambda = q \log M - p \log 2.$$

It is nonzero: otherwise $M^q = 2^p$, which would make $\beta = p/q$ rational; a rational solution is integral by the elementary rational-exponent classification already formalized in the source corpus. The Baker–Wüstholz theorem, or Matveev’s explicit lower bound for a homogeneous linear form in logarithms of algebraic numbers, gives

$$\log |\Lambda| \geq -C(1 + \log M) \log H - C'(1 + \log M),$$

where $H \geq \max(3, |p|, q)$ and C, C' are effective absolute constants after the fixed number 2 is absorbed. For approximations with $|p| \leq (\beta + 1)q$, one has $H \ll_M q$; for all other p the desired inequality is trivial after decreasing c_M . Since

$$\left| \beta - \frac{p}{q} \right| = \frac{|\Lambda|}{q \log 2},$$

absorbing fixed factors and one additional power of q gives (110). Matveev’s dependence on the logarithmic heights is linear in $\log M = \beta \log 2$ here, which gives the final scale assertion. $\square$

Status: [classical/open background] plus [paper proof in this report]

The lower-bound theorem for logarithmic forms is classical [3, 19, 20]. Its application to the primitive generator, the continued-fraction consequences below, and the comparison with the source report are new deductions in this continuation. No useful small numerical value of μ_M is claimed; standard explicit constants are very large.

34.2 Polynomial recurrence for convergent denominators

Let p_n/q_n be the continued-fraction convergents of β . The standard upper bound

$$\left| \beta - \frac{p_n}{q_n} \right| < \frac{1}{q_n q_{n+1}}$$

combined with (110) gives the following.

Corollary 34.2 (Polynomial denominator growth). *There is an effective constant $C_M > 0$ such that*

$$q_{n+1} \leq C_M q_n^{\mu_M - 1} \quad (111)$$

for every n with $q_n \geq 2$. Consequently the next partial quotient satisfies $a_{n+1} = O_M(q_n^{\mu_M - 2})$.

This is qualitatively stronger than the inherited elementary estimate $q_{n+1} < 2M^{q_n}$. It rules out Liouville behavior and says that the exceptionally long near-repetitions of the Sturmian word are polynomially controlled in the preceding standard word length. The exponent depends on the unknown integer M , however, and may be enormous.

34.3 Polynomial spacing in conditional blocks

The irrationality measure also upgrades the separation of the block nodes. From (110), for $d \geq 1$ and the nearest integer r to $d\beta$,

$$\|d\beta\|_{\mathbb{Z}} = |d\beta - r| \geq c_M d^{1-\mu_M}. \quad (112)$$

For an integer $T \geq 0$, put

$$K_T = \left\lfloor \frac{T+1}{\beta} \right\rfloor, \quad x_{T,k} = T - \lfloor k\beta \rfloor + k\beta = T + \{k\beta\}, \quad 0 \leq k \leq K_T.$$

The associated candidate and output nodes are

$$m_{T,k} = 2^{x_{T,k}} = 2^{T-\lfloor k\beta \rfloor} M^k, \quad a_{T,k} = 3^{x_{T,k}} = 3^{T-\lfloor k\beta \rfloor} A^k. \quad (113)$$

Corollary 34.3 (Effective block separation). *For distinct $i, j \in \{0, \dots, K_T\}$,*

$$|x_{T,j} - x_{T,i}| \geq c_M K_T^{1-\mu_M}, \quad (114)$$

$$|m_{T,j} - m_{T,i}| \geq (\log 2) c_M 2^T K_T^{1-\mu_M}, \quad (115)$$

$$|a_{T,j} - a_{T,i}| \geq (\log 3) c_M 3^T K_T^{1-\mu_M}. \quad (116)$$

Proof. The ordinary distance between two fractional parts is at least their circular distance, so

$$|\{j\beta\} - \{i\beta\}| \geq \|(j-i)\beta\|_{\mathbb{Z}} \geq c_M |j-i|^{1-\mu_M} \geq c_M K_T^{1-\mu_M},$$

because $1 - \mu_M < 0$. This proves (114). The mean value theorem applied to 2^x and 3^x on $[T, T+1]$ gives the remaining inequalities. $\square$

34.4 Separation on the full (n, k) lattice

The same estimate applies to the two-dimensional conditional exponent grid.

Corollary 34.4 (Polynomial separation of conditional exponents). *There are effective $c_M, \mu_M > 0$ such that, for distinct integer pairs $(n, k), (n', k')$ with $0 \leq k, k' \leq H$,*

$$|(n + k\beta) - (n' + k'\beta)| \geq \min\{1, c_M H^{1-\mu_M}\}.$$

Proof. If $k = k'$, the difference is a nonzero integer. Otherwise divide the linear form $(n - n') + (k - k')\beta$ by $|k - k'|$ and apply (110). $\square$

The column exponent lattice $u + \vartheta v$ has an analogous effective polynomial spacing with absolute constants, because $\vartheta = \log_2 3$ is a fixed ratio of two logarithms. More refined explicit estimates are known for linear forms in $\log 2$ and $\log 3$; for example Wu’s simultaneous linear-independence measure leads, in the usual normalization, to an irrationality exponent for ϑ below approximately 8.62 [24]. The exact convention-dependent numerical conversion is not needed below.

34.5 Why the Baker upgrade does not close the conjecture

The effective spacing is useful, but its scale is still compatible with the exact conditional population. A dyadic exponent block contains only $K_T + 1 \asymp T$ nodes. A polynomial lower bound on their minimum spacing imposes no contradiction on T points in a fixed interval. Likewise, the full triangular grid has $\asymp T^2$ points in an interval of exponent length T , and polynomially small gaps are expected.

Baker’s theorem therefore strengthens compactness and determinant estimates but does not supply the missing logarithm by itself. Its most promising roles are indirect:

- controlling the smallest row and column gaps in an exact Harish–Chandra–Itzykson–Zuber factorization;
- bounding how often a p -adic assignment problem can be nearly degenerate;
- limiting the lengths of near-periodic Sturmian blocks that a congruence argument must handle;
- making computational searches over continued-fraction branches provably finite once a candidate output M is fixed.

35 Exact all-support valuations of ordinary block Vandermondes

35.1 The block determinants

Fix an integer T and abbreviate $K = K_T$. Let $m_k = m_{T,k}$ and $a_k = a_{T,k}$ be the $K + 1$ input–output pairs in (113). Their ordinary Vandermonde determinants are

$$V_T^{(2)} = \det(m_i^j)_{0 \leq i, j \leq K} = \prod_{0 \leq i < j \leq K} (m_j - m_i), \quad (117)$$

$$V_T^{(3)} = \det(a_i^j)_{0 \leq i, j \leq K} = \prod_{0 \leq i < j \leq K} (a_j - a_i). \quad (118)$$

They are nonzero integers. Since $m_k \in [2^T, 2^{T+1})$ and $a_k \in [3^T, 3^{T+1})$,

$$\log |V_T^{(2)}| \leq R_T T \log 2, \quad (119)$$

$$\log |V_T^{(3)}| \leq R_T (T \log 3 + \log 2), \quad (120)$$

where

$$R_T = \binom{K+1}{2} = \frac{K(K+1)}{2}.$$

The additive $R_T \log 2$ in (120) is lower order.

35.2 The input determinant: an exact formula at every support prime

Factor

$$M = 2^a W, \quad W \text{ odd}, \quad \delta_2 = \beta - a = \log_2 W.$$

As noted above, $W > 1$, so $W \geq 3$ and

$$\delta_2 \geq \log_2 3 > 1. \quad (121)$$

Because a is an integer,

$$m_k = 2^{T - \lfloor k\delta_2 \rfloor} W^k. \quad (122)$$

The 2-adic valuations in (122) strictly decrease with k , while every odd support-prime valuation strictly increases.

Proposition 35.1 (Exact support valuations for $V_T^{(2)}$). *For $0 \leq i < j \leq K$,*

$$v_2(m_j - m_i) = T - \lfloor j\delta_2 \rfloor, \quad (123)$$

$$v_p(m_j - m_i) = i v_p(W) \quad (p \mid W). \quad (124)$$

Consequently

$$v_2(V_T^{(2)}) = \sum_{j=1}^K j(T - \lfloor j\delta_2 \rfloor), \quad (125)$$

$$v_p(V_T^{(2)}) = v_p(W) \sum_{i=0}^{K-1} i(K-i) \quad (p \mid W). \quad (126)$$

Proof. Put $\rho_k = T - \lfloor k\delta_2 \rfloor$. By (121), $\rho_i > \rho_j$ when $i < j$. From (122),

$$m_j - m_i = 2^{\rho_j} W^i \left(W^{j-i} - 2^{\rho_i - \rho_j} \right).$$

The parenthesis is odd, proving (123). For an odd prime $p \mid W$, the two terms m_i and m_j have distinct valuations $i v_p(W) < j v_p(W)$, so the ultrametric equality gives (124). Summing over pairs gives (125)–(126). $\square$

Let $Q_T^{(2)}$ be the part of $V_T^{(2)}$ supported on the primes dividing $2M$. Then Proposition 35.1 gives the exact identity

$$\log Q_T^{(2)} = \log 2 \sum_{j=1}^K j(T - \lfloor j\delta_2 \rfloor) + \log W \sum_{i=0}^{K-1} i(K-i). \quad (127)$$

The two elementary sums have asymptotics

$$\begin{aligned} \sum_{j=1}^K j(T - \lfloor j\delta_2 \rfloor) &= \frac{TK^2}{2} - \frac{\delta_2 K^3}{3} + O_\beta(T^2), \\ \sum_{i=0}^{K-1} i(K - i) &= \frac{K^3}{6} + O(K^2). \end{aligned}$$

Using $\log W = \delta_2 \log 2$ and $K/T \rightarrow 1/\beta$ yields

$$\frac{\log Q_T^{(2)}}{R_T T \log 2} \rightarrow 1 - \frac{\delta_2}{3\beta}. \quad (128)$$

This includes every prime appearing in the entries. Any further prime divisor of the Vandermonde comes from subtraction and hence from genuine cancellation beyond termwise entry valuations.

35.3 The output determinant and bounded tie corrections

Factor

$$A = 3^b Z, \quad 3 \nmid Z, \quad \delta_3 = \beta - b = \log_3 Z > 0.$$

Then

$$a_k = 3^{T - \lfloor k\delta_3 \rfloor} Z^k. \quad (129)$$

Unlike δ_2 , the number δ_3 can be smaller than 1; for example the least possible base-free factor is 2. Equal consecutive 3-adic valuations can therefore occur. They contribute only a lower-order correction.

Proposition 35.2 (Output support valuations). *For every prime $p \mid Z$ and $0 \leq i < j \leq K$,*

$$v_p(a_j - a_i) = i v_p(Z).$$

At $p = 3$,

$$v_3(a_j - a_i) = T - \lfloor j\delta_3 \rfloor$$

after adding a nonnegative correction that is zero whenever $\lfloor i\delta_3 \rfloor < \lfloor j\delta_3 \rfloor$. The total correction over all pairs is $O_{A,\beta}(K^2)$. Consequently, if $Q_T^{(3)}$ denotes the part of $V_T^{(3)}$ supported on primes dividing $3A$, then

$$\frac{\log Q_T^{(3)}}{R_T T \log 3} \rightarrow 1 - \frac{\delta_3}{3\beta}. \quad (130)$$

Proof. The assertion for $p \mid Z$ is the same distinct-valuation argument as before. Put $\rho_k = T - \lfloor k\delta_3 \rfloor$. If $\rho_i > \rho_j$, then

$$a_j - a_i = 3^{\rho_j} Z^i \left(Z^{j-i} - 3^{\rho_i - \rho_j} \right)$$

and the parenthesis is not divisible by 3, so the valuation is exactly ρ_j . If $\rho_i = \rho_j$, then

$$v_3(a_j - a_i) = \rho_j + v_3(Z^{j-i} - 1).$$

Equality of the floors implies $(j - i)\delta_3 < 1$, so $j - i$ belongs to a fixed finite set depending only on δ_3 . Hence the extra valuation is bounded by a constant depending on A and β . There are $O(K^2)$ pairs. The asymptotic calculation is then identical to (127). $\square$

35.4 The depth-sensitive product ceiling

Combining (119), (120), (128), and (130) gives the principal conclusion of the ordinary Vandermonde audit.

Theorem 35.3 (All-support ordinary Vandermonde fraction). *For a hypothetical primitive counterexample,*

$$\frac{\log(Q_T^{(2)} Q_T^{(3)})}{R_T T \log 6} \longrightarrow 1 - \frac{\delta_2 \log 2 + \delta_3 \log 3}{3\beta \log 6}, \quad \delta_2 = \beta - a, \quad \delta_3 = \beta - b. \quad (131)$$

Because $\min(a, b) = 0$, the limit is at most

$$1 - \frac{\log 2}{3 \log 6} = 0.871049\,064\dots \quad (132)$$

If $b = 0$, the sharper ceiling is

$$1 - \frac{\log 3}{3 \log 6} = 0.795599\,921\dots$$

If $a = b = 0$, the limit is exactly $2/3$.

Proof. Only the numerical ceiling remains to be shown. If $a = 0$, then $\delta_2 = \beta$, so the deficit in (131) is at least $\log 2/(3 \log 6)$. If $b = 0$, it is at least $\log 3/(3 \log 6)$. Taking the larger possible fraction gives (132). If both vanish, the numerator of the deficit is $\beta \log 6$. $\square$

Table 6: Leading termwise-divisor ceilings for ordinary block Vandermondes.

Branch information	Exact/upper fraction	Decimal
$a = b = 0$	$2/3$	$0.666666\dots$
$b = 0$	$\leq 1 - \log 3/(3 \log 6)$	$0.795599\dots$
No matched depth only	$\leq 1 - \log 2/(3 \log 6)$	$0.871049\dots$

35.5 Why even 87% is not enough

The deficit in Theorem 35.3 is a fixed positive fraction once the hypothetical primitive pair is fixed. The archimedean height is $\asymp T^3$, whereas all factorial, spacing, and tie corrections are $O_\beta(T^2 \log T)$. Thus no lower-order refinement of the ordinary Vandermonde can close the remaining gap.

This conclusion is stronger than the earlier observation that the 2-adic valuation alone captures one third in the base-free case. It includes every prime already present in the entries and every possible base depth. A successful ordinary-Vandermonde proof would need a new leading-order source of divisibility in the prime factors of the *differences*; the fixed-support valuations are completely accounted for by Theorem 35.3.

36 The mixed-semigroup determinant and its Gini norm

36.1 Balanced low-weight columns

Let

$$\Lambda = \{u + \vartheta v : (u, v) \in \mathbb{N}_0^2\}.$$

Because ϑ is irrational, every element has a unique pair (u, v) . List the pairs in increasing order of $u + \vartheta v$:

$$(u_0, v_0), (u_1, v_1), \dots$$

For $N \geq 1$, let

$$L_N = \max_{0 \leq j < N} (u_j + \vartheta v_j).$$

The elementary lattice count

$$\#\{(u, v) \in \mathbb{N}_0^2 : u + \vartheta v \leq L\} = \frac{L^2}{2\vartheta} + O_\vartheta(L) \quad (133)$$

shows that

$$L_N \sim \sqrt{2\vartheta N}. \quad (134)$$

Moreover, the empirical distribution of

$$X_j = \left(\frac{u_j}{L_N}, \frac{\vartheta v_j}{L_N} \right) \quad (135)$$

converges to uniform area measure on the standard triangle

$$\mathcal{T} = \{(x, z) \in \mathbb{R}^2 : x \geq 0, z \geq 0, x + z \leq 1\}. \quad (136)$$

Here and below “uniform” means probability measure with density 2 on this triangle.

For the block at height T , set $K = K_T$ and $N = K + 1$. Define

$$D_T = \det \left(m_{T,k}^{u_j} a_{T,k}^{v_j} \right)_{0 \leq k, j < N}. \quad (137)$$

The generalized-Vandermonde theorem inherited from the source report implies that D_T is a nonzero integer. Put

$$B_j = 2^{u_j} 3^{v_j}.$$

Since $m_{T,k} = 2^{x_{T,k}}$ and $a_{T,k} = 3^{x_{T,k}}$,

$$m_{T,k}^{u_j} a_{T,k}^{v_j} = B_j^{x_{T,k}} = B_j^{T + \{k\beta\}}. \quad (138)$$

Thus the Leibniz expansion gives the clean archimedean bound

$$\log |D_T| \leq (T+1) \sum_{j < N} \log B_j + \log N!. \quad (139)$$

The leading height is

$$H_T = T \sum_{j < N} \log B_j. \quad (140)$$

By (134) and the mean $\mathbb{E}(X_1 + X_2) = 2/3$ on $\mathcal{T}$,

$$H_T \sim \frac{2}{3} T N L_N \log 2 \asymp_\beta T^{5/2}. \quad (141)$$

Both $\sum_j \log B_j$ and $\log N!$ in the error of (139) are $o(H_T)$.

36.2 The all-prime tropical divisor

For a prime p , define the assignment minimum

$$\tau_{p,T} = \min_{\sigma \in S_N} \sum_{k=0}^{N-1} v_p \left(m_{T,k}^{u_{\sigma(k)}} a_{T,k}^{v_{\sigma(k)}} \right). \quad (142)$$

Every Leibniz term is divisible by $p^{\tau_{p,T}}$, so

$$Q_T^{\text{trop}} = \prod_p p^{\tau_{p,T}} \quad \text{divides} \quad D_T. \quad (143)$$

Only primes dividing $6MA$ occur in the product.

The problem is now to compute the leading asymptotic of $\log Q_T^{\text{trop}}/H_T$. A crude estimate “minimum $\leq$ average” loses the precise branch dependence. The exact answer is controlled by a simple probability norm.

36.3 Antimonotone assignment and the Gini norm

Let $X = (X_1, X_2)$ be uniform on $\mathcal{T}$, and let X' be an independent copy. For $c = (r, s) \in \mathbb{R}^2$, define

$$\mathfrak{G}(r, s) = \mathbb{E} |r(X_1 - X'_1) + s(X_2 - X'_2)|. \quad (144)$$

This is a norm on $\mathbb{R}^2$: homogeneity and the triangle inequality are immediate, and the support of $X - X'$ spans $\mathbb{R}^2$, so only $(r, s) = (0, 0)$ has norm zero.

Let $Y = rX_1 + sX_2$, with ascending quantile $Q_Y(t)$. Pairing an increasing uniform row parameter with the decreasing ordering of Y gives the continuum assignment functional

$$\Phi(r, s) = \int_0^1 t Q_Y(1-t) dt. \quad (145)$$

Lemma 36.1 (Gini formula for the minimum assignment). *For every $(r, s) \in \mathbb{R}^2$,*

$$\Phi(r, s) = \frac{1}{2} \mathbb{E} Y - \frac{1}{4} \mathfrak{G}(r, s) = \frac{r+s}{6} - \frac{1}{4} \mathfrak{G}(r, s). \quad (146)$$

Proof. The rearrangement inequality says that the minimum scalar product pairs the increasing row parameter with the decreasing list of column values. Changing variables $u = 1 - t$ in (145) gives

$$\Phi = \int_0^1 (1-u) Q_Y(u) du.$$

If Y' is an independent copy, the minimum of Y, Y' has quantile density $2(1-u)$, so

$$\Phi = \frac{1}{2} \mathbb{E} \min(Y, Y').$$

Since $\min(Y, Y') = (Y + Y' - |Y - Y'|)/2$, this becomes $\frac{1}{2} \mathbb{E} Y - \frac{1}{4} \mathbb{E} |Y - Y'|$. Finally $\mathbb{E} X_1 = \mathbb{E} X_2 = 1/3$. $\square$

36.4 Closed form of the Gini norm

The norm (144) is elementary. Sort the three vertex values

$$0, \quad r, \quad s$$

as $\xi_0 \leq \xi_1 \leq \xi_2$, and put

$$A_0 = \xi_1 - \xi_0, \quad B_0 = \xi_2 - \xi_1.$$

Proposition 36.2 (Explicit triangular Gini norm). *If $A_0 + B_0 > 0$, then*

$$\mathfrak{G}(r, s) = \frac{2(2A_0^2 + 3A_0B_0 + 2B_0^2)}{15(A_0 + B_0)}. \quad (147)$$

At $(r, s) = (0, 0)$ the value is 0. In particular,

$$\mathfrak{G}(1, 0) = \mathfrak{G}(0, 1) = \mathfrak{G}(1, 1) = \frac{4}{15}, \quad \mathfrak{G}(1, -1) = \frac{7}{15}. \quad (148)$$

Proof. Translation of the three vertex values does not change a difference $Y - Y'$, so take $\xi_0 = 0$, $\xi_1 = A_0$, and $\xi_2 = A_0 + B_0 = L$. Projection of uniform area measure on a triangle by a linear function has density

$$f(y) = \begin{cases} \frac{2y}{A_0L}, & 0 \leq y \leq A_0, \\ \frac{2(L-y)}{B_0L}, & A_0 \leq y \leq L. \end{cases}$$

Thus its distribution function is $F(y) = y^2/(A_0L)$ on the first interval and $F(y) = 1 - (L-y)^2/(B_0L)$ on the second. The standard identity

$$\mathbb{E}|Y - Y'| = 2 \int_{-\infty}^{\infty} F(y)(1 - F(y)) dy$$

gives (147) after two polynomial integrations. Cases with one gap zero follow by continuity. $\square$

For later use, if $\gamma \geq 0$, then

$$\mathfrak{G}(-1, \gamma) = \frac{2(2 + 3\gamma + 2\gamma^2)}{15(1 + \gamma)} = \frac{4}{15} + \frac{2\gamma(1 + 2\gamma)}{15(1 + \gamma)} \geq \frac{4}{15}. \quad (149)$$

36.5 Normalized prime-valuation slope vectors

Recall

$$a = v_2(M), \quad b = v_3(A), \quad c = v_2(A), \quad d = v_3(M).$$

For each prime, define a normalized slope vector $c_p \in \mathbb{R}^2$ by

$$c_2 = \left(-1 + \frac{a}{\beta}, \frac{c}{\beta\vartheta} \right), \quad (150)$$

$$c_3 = \left(\frac{\vartheta d}{\beta}, -1 + \frac{b}{\beta} \right), \quad (151)$$

$$c_p = \left(\frac{v_p(M) \log p}{\beta \log 2}, \frac{v_p(A) \log p}{\beta \log 3} \right), \quad p \neq 2, 3. \quad (152)$$

The logarithmic factorizations of M and A give the conservation law

$$\sum_p c_p = (0, 0). \quad (153)$$

Indeed, the first coordinate sum is

$$-1 + \frac{a \log 2 + d \log 3 + \sum_{p \neq 2,3} v_p(M) \log p}{\beta \log 2} = 0,$$

and similarly for the second coordinate.

The meaning of these vectors can be read directly from the valuation matrices. Put $s_k = \beta k/T$. Up to $o(L_N)$ after division by T , the logarithmic p -adic cost of the entry indexed by a scaled column (x, z) is

Prime	Cost divided by $TL_N \log 2$
2	$x + s_k c_2 \cdot (x, z)$
3	$z + s_k c_3 \cdot (x, z)$
$p \neq 2, 3$	$s_k c_p \cdot (x, z)$

For example,

$$v_2(m_k^u a_k^v) = u(T - \lfloor k\beta \rfloor + ak) + vck,$$

and substituting $u = L_N x$, $v = L_N z/\vartheta$, $k \sim Ts_k/\beta$ gives the first row. The $O(u)$ errors from the floors sum to $O(NL_N) = o(H_T)$.

36.6 The exact continuum limit

Theorem 36.3 (All-prime tropical fraction for the mixed block determinant). *Under a hypothetical primitive counterexample, the determinant (137) satisfies*

$$\frac{\log Q_T^{\text{trop}}}{H_T} \longrightarrow 1 - \frac{3}{8} \sum_p \mathfrak{G}(c_p), \quad (154)$$

where the finitely many slope vectors c_p are given by (150)–(152).

Proof. The row parameters $s_k = \beta k/T$ become uniformly distributed on $[0, 1]$, because $K_T \sim T/\beta$. The scaled columns become uniformly distributed on $\mathcal{T}$ by (133). The finite rearrangement inequality and Riemann-sum convergence therefore show that the assignment minimum at prime p , after normalization by $TL_N N \log 2$, tends to its constant-column mean plus $\Phi(c_p)$.

Only $p = 2$ has constant part x , and only $p = 3$ has constant part z . Their combined mean is $2/3$. By Lemma 36.1, the sum of the slope contributions is

$$\sum_p \Phi(c_p) = \frac{1}{6} \sum_p (c_{p,1} + c_{p,2}) - \frac{1}{4} \sum_p \mathfrak{G}(c_p).$$

The first sum vanishes by (153). Hence

$$\frac{\log Q_T^{\text{trop}}}{TL_N N \log 2} \longrightarrow \frac{2}{3} - \frac{1}{4} \sum_p \mathfrak{G}(c_p).$$

On the other hand, (141) gives $H_T/(TL_N N \log 2) \rightarrow 2/3$. Dividing proves (154). $\square$

This theorem is exact at the level of minimum Leibniz-term valuations. It replaces a family of separate p -adic estimates by one finite norm sum satisfying a conservation law.

The four base and cross depths already give a useful branch-sensitive bound without knowing any external prime factorization.

Corollary 36.4 (Depth-sensitive three-vector bound). *For every hypothetical primitive counterexample,*

$$\limsup_{T \rightarrow \infty} \frac{\log Q_T^{\text{trop}}}{H_T} \leq 1 - \frac{3}{8} \left(\mathfrak{G}(c_2) + \mathfrak{G}(c_3) + \mathfrak{G}(-c_2 - c_3) \right). \quad (155)$$

Equality in this upper bound can occur only when all nonzero external slope vectors are positively collinear.

Proof. The external vectors satisfy $\sum_{p \neq 2,3} c_p = -c_2 - c_3$. Hence the triangle inequality gives

$$\sum_{p \neq 2,3} \mathfrak{G}(c_p) \geq \mathfrak{G}(-c_2 - c_3).$$

Substitute this in Theorem 36.3. Equality in a norm triangle inequality requires all summands to lie on one nonnegative ray. $\square$

The equality condition detects exactly the second-iterate kernel branch reviewed in the source report. Indeed the external vectors c_p are rescaled versions of $(v_p(M), v_p(A))$. They are all collinear if and only if the two external valuation vectors of M and A are rationally dependent, equivalently if the kernel K_β is nonzero. Thus:

Corollary 36.5 (The tropical norm sees the kernel dichotomy). *In the multiplicatively independent branch $K_\beta = 0$, the inequality in (155) is strict. In the dependent branch $K_\beta \neq 0$, the external triangle inequality is an equality (including the axis cases in which one output is $\{2, 3\}$ -smooth). For fixed M, A , the strict gap in the independent branch is explicit:*

$$\Delta_{\text{ext}} = \sum_{p \neq 2,3} \mathfrak{G}(c_p) - \mathfrak{G} \left(\sum_{p \neq 2,3} c_p \right) > 0. \quad (156)$$

There is no uniform positive lower bound for Δ_{ext} : distinct rational rays can approach one another as the valuations grow. Nevertheless, the quantity is a computable branch invariant and gives a sharper determinant threshold for every fixed hypothetical primitive pair.

36.7 The universal 4/5 ceiling

Theorem 36.6 (Sharp universal min-plus ceiling). *For every hypothetical primitive counterexample,*

$$\sum_p \mathfrak{G}(c_p) \geq \frac{8}{15}. \quad (157)$$

Consequently

$$\limsup_{T \rightarrow \infty} \frac{\log Q_T^{\text{trop}}}{H_T} \leq \frac{4}{5}. \quad (158)$$

The constant 4/5 is the sharp supremum allowed by the normalized valuation constraints, although equality would require a degenerate boundary incompatible with a genuine nonintegral solution.

Proof. By the primitive theorem, either $a = 0$ or $b = 0$. Suppose first that $a = 0$. Then $c_2 = (-1, \gamma)$ with $\gamma = c/(\beta\vartheta) \geq 0$. By (149), $\mathfrak{G}(c_2) \geq 4/15$. Conservation and the triangle inequality for the norm give

$$\sum_{p \neq 2} \mathfrak{G}(c_p) \geq \mathfrak{G}\left(\sum_{p \neq 2} c_p\right) = \mathfrak{G}(-c_2) = \mathfrak{G}(c_2).$$

Thus the total is at least $8/15$. If $b = 0$, apply the same argument to $c_3 = (\eta, -1)$, where $\eta = \vartheta d/\beta \geq 0$. Substitution in (154) gives $1 - (3/8)(8/15) = 4/5$. $\square$

The proof is short because the difficult arithmetic has been compressed into two facts: the slope vectors sum to zero, and one structural vector reaches from -1 to a nonnegative vertex. No averaging estimate alone reveals this geometry.

36.8 Cross-coprime constants: $7/10$ and $3/10$

Assume the stronger branch condition

$$\gcd(MA, 6) = 1. \quad (159)$$

Then $a = b = c = d = 0$,

$$c_2 = (-1, 0), \quad c_3 = (0, -1), \quad \sum_{p \neq 2, 3} c_p = (1, 1).$$

Corollary 36.7 (All-prime cross-coprime ceiling). *Under (159),*

$$\limsup_{T \rightarrow \infty} \frac{\log Q_T^{\text{trop}}}{H_T} \leq \frac{7}{10}. \quad (160)$$

If the external normalized valuation vectors are not all positively collinear, the inequality is strict by a branch-dependent positive amount.

Proof. By (148), the structural vectors each contribute $4/15$. The triangle inequality gives

$$\sum_{p \neq 2, 3} \mathfrak{G}(c_p) \geq \mathfrak{G}(1, 1) = \frac{4}{15}.$$

Thus the total norm is at least $4/5$, and $1 - (3/8)(4/5) = 7/10$. Equality in the triangle inequality requires positive collinearity. $\square$

The places 2 and 3 alone have an even smaller exact continuum contribution.

Corollary 36.8 (The structural-place $3/10$ constant). *Under (159), let $Q_T^{(2,3)} = 2^{\tau_2, T} 3^{\tau_3, T}$. Then*

$$\frac{\log Q_T^{(2,3)}}{H_T} \longrightarrow \frac{3}{10}. \quad (161)$$

Proof. For $c = (-1, 0)$,

$$\Phi(c) = \frac{-1}{6} - \frac{1}{4} \cdot \frac{4}{15} = -\frac{7}{30},$$

and the same holds for $(0, -1)$. Adding the constant mean $2/3$ gives $2/3 - 7/15 = 1/5$. Dividing by the height mean $2/3$ gives $3/10$. $\square$

Equivalently, one can compute this constant from quantiles. The descending quantile of X_1 on $\mathcal{T}$ is $1 - \sqrt{t}$, so

$$\int_0^1 t(1 - \sqrt{t}) dt = \frac{1}{10}, \quad \mathbb{E}X_1 = \frac{1}{3}.$$

The ratio is $(1/10)/(1/3) = 3/10$ in each structural coordinate.

36.9 Correction and calibration of a later repository note

A short p -adic priority note added to the repository after the attached unified report makes the correct strategic observation that minimum Leibniz-term valuations leave a constant fraction of the archimedean height uncovered. Two details require correction.

First, the displayed structural valuation formulas in that note simultaneously omit the terms $kuv_2(M)$ and $kvv_3(A)$; they therefore assume both $a = 0$ and $b = 0$. The inherited primitive theorem gives only $\min(a, b) = 0$. Equations (150)–(151) retain all four cross/base depths.

Second, the proposed all-prime $2/3$ ceiling is not a valid universal consequence of “minimum $\leq$ average.” For an external positive weight proportional to $X_1 + X_2$, the minimum assignment is $4/5$ of its average assignment, not at most $2/3$. The exact Gini formula shows that the universal normalized ceiling is $4/5$, while the cross-coprime all-prime ceiling is $7/10$. The note’s qualitative conclusion survives—entrywise valuation is insufficient—but the constants and branch hypotheses are materially different.

Status: [\[paper proof in this report\]](#)

The correction is not based on a numerical counterexample to the Alaoglu–Erdős conjecture. It is an exact audit of the assignment inequality used in the method. Rational external valuation directions can approximate the equality configurations arbitrarily well, so no stronger uniform constant follows merely from integrality of the valuation vectors.

36.10 Where extra divisibility can occur

The tropical minimum equals the true determinant valuation whenever its minimizing permutation is unique.

Proposition 36.9 (Unique tropical term forbids cancellation). *Let $D = \det(E_{ij})$ be a nonzero integer determinant. Fix a prime p , and suppose a unique permutation σ_0 minimizes $\sum_i v_p(E_{i, \sigma_0(i)})$. Then*

$$v_p(D) = \sum_i v_p(E_{i, \sigma_0(i)}).$$

Proof. Divide the Leibniz expansion by the corresponding power of p . The term indexed by σ_0 is a p -adic unit, while every other term is divisible by p . The sum is therefore a unit modulo p . $\square$

Consequently, the missing 20% or more can arise only from primes at which the assignment problem has several minimizers. The tie geometry is explicit.

- For $p \neq 2, 3$, columns tie when $v_p(M)(u - u') + v_p(A)(v - v') = 0$, so ties lie on rational lines orthogonal to the external valuation vector.

- In the cross-coprime branch, the 2-adic cost depends only on u and the 3-adic cost only on v . Every vertical or horizontal column fiber creates many equal-cost permutations. Any leading extra valuation must be produced inside these fibers.
- Cross valuations $v_2(A)$ and $v_3(M)$ tilt the fibers. The floor term in the row profile perturbs the exact affine ordering but does not change the continuum slope vector.

This turns “look for p -adic cancellation” into a concrete algebra problem: factor or control the residual determinants inside equal-slope column fibers, and show that their valuation supplies at least the missing Gini deficit.

37 The complete triangular grid and a $\pi/4$ transport ceiling

The one-block determinant deliberately uses the exact $O(T)$ population in a unit interval of exponent space. The source report also suggests exploiting the full two-dimensional coordinate grid rather than sorting only by magnitude. Conditional on a counterexample, that proposal gives $\asymp T^2$ genuine integer rows below exponent height T , exactly the number that the mixed-semigroup interpolation argument seemed to be missing. It is therefore important to determine whether entrywise p -adic divisibility becomes strong enough on this larger node set.

The answer is again negative, but for a different and pleasantly universal reason. Once all places are summed, the prime-by-prime assignment minima are bounded by one antimonotone transport problem for the logarithmic entry matrix. Both the row heights and the balanced column weights have the same triangular distribution, and the ratio between antimonotone and monotone transport is exactly $\pi/4$.

37.1 Rows and columns

For $T > 0$ define the complete conditional row set

$$\mathcal{R}_T = \{(n, k) \in \mathbb{N}_0^2 : n + k\beta \leq T\}, \quad N_T = \#\mathcal{R}_T. \quad (162)$$

For $i = (n_i, k_i) \in \mathcal{R}_T$ put

$$x_i = n_i + k_i\beta, \quad m_i = 2^{n_i} M^{k_i} = 2^{x_i}, \quad a_i = 3^{n_i} A^{k_i} = 3^{x_i}. \quad (163)$$

Because β is irrational, the x_i are pairwise distinct.

Let $\mathcal{C}_T$ be the first N_T points $(u, v) \in \mathbb{N}_0^2$ in increasing order of

$$\lambda = u + \vartheta v, \quad y = \log(2^u 3^v) = \lambda \log 2.$$

Write L_T for the largest selected λ . Standard lattice-point estimates for the two triangles give

$$N_T = \frac{T^2}{2\beta} + O_\beta(T), \quad L_T = \sqrt{2\vartheta N_T} + O_\vartheta(1) = T \sqrt{\frac{\vartheta}{\beta}} + O_\beta(1). \quad (164)$$

The associated square determinant is

$$D_T^\Delta = \det \left(m_i^{u_j} a_i^{v_j} \right)_{i,j=1}^{N_T} = \det (e^{x_i y_j})_{i,j=1}^{N_T}. \quad (165)$$

Every entry is an integer. Ordering the x_i and y_j increasingly, strict total positivity of the exponential kernel makes $D_T^\Delta > 0$; equivalently, this is the same exponential-polynomial zero theorem used by the generalized-Vandermonde part of the source report.

Status: [\[paper proof in this report\]](#) for the asymptotic transport statements below. Integrality and nonvanishing are inherited from the mixed-semigroup determinant architecture [\[kernel-verified Lean\]](#).

37.2 The common triangular distribution

The normalized row heights $h_i = x_i/T$ and column weights $\ell_j = y_j/(L_T \log 2)$ both live in $[0, 1]$. Their empirical measures have the same limit.

Lemma 37.1 (Triangular height law). *For every continuous $f : [0, 1] \rightarrow \mathbb{R}$,*

$$\frac{1}{N_T} \sum_{i=1}^{N_T} f(h_i) \longrightarrow \int_0^1 2t f(t) dt, \quad (166)$$

$$\frac{1}{N_T} \sum_{j=1}^{N_T} f(\ell_j) \longrightarrow \int_0^1 2t f(t) dt. \quad (167)$$

Thus the common limiting distribution has cumulative distribution function $F(t) = t^2$ and quantile function

$$Q(s) = \sqrt{s}, \quad 0 \leq s \leq 1. \quad (168)$$

Proof. After scaling (n, k) to $(n/T, \beta k/T)$, the row lattice becomes equidistributed in the unit right triangle

$$\{(r, s) : r, s \geq 0, r + s \leq 1\}.$$

The subtriangle on which $r + s \leq t$ has area $t^2/2$, while the whole triangle has area $1/2$. Hence the height $r + s$ has distribution function t^2 . Boundary discrepancies are $O(T)$ against $N_T \asymp T^2$, proving (166).

For the columns, scale (u, v) to $(u/L_T, \vartheta v/L_T)$. The same unit triangle appears, and $\ell = (u + \vartheta v)/L_T$ is again the sum of the scaled coordinates. The boundary shell created by taking the first exactly N_T points has $O(L_T)$ points against $N_T \asymp L_T^2$, which proves (167). $\square$

37.3 Summing all primewise assignment minima

For each prime p , let

$$\tau_{p,T} = \min_{\sigma \in S_{N_T}} \sum_{i=1}^{N_T} v_p \left(m_i^{u_{\sigma(i)}} a_i^{v_{\sigma(i)}} \right), \quad Q_T^\Delta = \prod_p p^{\tau_{p,T}}. \quad (169)$$

As before, Q_T^Δ divides every Leibniz term and hence divides D_T^Δ . The key observation is that taking a separate minimizing permutation at each prime can only decrease the sum relative to using one permutation at all primes.

Lemma 37.2 (All-place transport domination). *For every T ,*

$$\log Q_T^\Delta \leq \min_{\sigma \in S_{N_T}} \sum_{i=1}^{N_T} x_i y_{\sigma(i)}. \quad (170)$$

Proof. For a permutation σ , put

$$C_p(\sigma) = \sum_i v_p \left(m_i^{u_{\sigma(i)}} a_i^{v_{\sigma(i)}} \right) \log p.$$

Then

$$\sum_p \min_{\sigma} C_p(\sigma) \leq \min_{\sigma} \sum_p C_p(\sigma).$$

Unique factorization and (163) give, entry by entry,

$$\sum_p v_p(m_i^{u_j} a_i^{v_j}) \log p = \log(m_i^{u_j} a_i^{v_j}) = x_i y_j.$$

Substitution proves (170). □

Notice how little arithmetic remains in the last inequality: it is valid in every valuation branch, includes all primes, and does not use the primitive-depth normalization. The price is that it is only an upper bound for the common termwise divisor. That is exactly what is needed for a no-go theorem.

37.4 Optimal transport and the $\pi/4$ ratio

Let

$$E_T^- = \min_{\sigma \in S_{N_T}} \sum_i x_i y_{\sigma(i)}, \quad (171)$$

$$E_T^+ = \max_{\sigma \in S_{N_T}} \sum_i x_i y_{\sigma(i)}. \quad (172)$$

The rearrangement inequality pairs the two ordered lists oppositely for E_T^- and in the same order for E_T^+ .

Theorem 37.3 (Triangular-grid transport constants). *As $T \rightarrow \infty$,*

$$\frac{E_T^-}{N_T T L_T \log 2} \longrightarrow \int_0^1 \sqrt{t(1-t)} dt = \frac{\pi}{8}, \quad (173)$$

$$\frac{E_T^+}{N_T T L_T \log 2} \longrightarrow \int_0^1 t dt = \frac{1}{2}. \quad (174)$$

Consequently

$$\frac{E_T^-}{E_T^+} \longrightarrow \frac{\pi}{4} = 0.7853981633 \dots \quad (175)$$

Proof. By Lemma 37.1, the increasing empirical quantile functions of the two normalized lists converge in L^1 to $Q(t) = \sqrt{t}$. The discrete rearrangement sums are Riemann sums for the monotone and antimonotone couplings. Thus

$$\frac{E_T^-}{N_T T L_T \log 2} \longrightarrow \int_0^1 Q(t)Q(1-t) dt, \quad \frac{E_T^+}{N_T T L_T \log 2} \longrightarrow \int_0^1 Q(t)^2 dt.$$

The first integral is the beta integral $B(3/2, 3/2) = \Gamma(3/2)^2/\Gamma(3) = \pi/8$; the second is $1/2$. $\square$

The determinant has $N_T!$ Leibniz terms, each at most $e^{E_T^+}$. Since $N_T \asymp T^2$ while $T L_T N_T \asymp T^4$,

$$\log(N_T!) = O(T^2 \log T) = o(N_T T L_T). \quad (176)$$

Combining the preceding statements gives the promised ceiling.

Corollary 37.4 (Full-grid termwise-divisor ceiling). *With the notation above,*

$$\limsup_{T \rightarrow \infty} \frac{\log Q_T^\Delta}{E_T^+ + \log(N_T!)} \leq \frac{\pi}{4}. \quad (177)$$

Thus the complete $\asymp T^2$ row grid does not make the all-prime common divisor exceed the sharp largest-Leibniz-term archimedean bound. Entrywise valuation alone leaves at least the fraction $1 - \pi/4 = 0.2146018\dots$ uncovered at leading order.

Proof. Lemma 37.2 gives $\log Q_T^\Delta \leq E_T^-$. Divide by the standard determinant upper bound $E_T^+ + \log(N_T!)$ and apply Theorem 37.3 and (176). $\square$

Caution 37.5 (What the $\pi/4$ theorem does not exclude). The denominator in (177) is the largest-term upper bound, not the true size of D_T^Δ . Generalized Vandermonde determinants exhibit substantial alternating cancellation. A sufficiently sharp archimedean estimate could lower the relevant budget below E_T^- , in which case the termwise divisor might become decisive. The theorem excludes the naive “more rows plus the same Hadamard/Leibniz estimate” route; it does not exclude a joint p -adic/free-energy argument.

37.5 The Harish–Chandra–Itzykson–Zuber identity

The exact analytic object behind that warning is classical. For diagonal Hermitian matrices $X = \text{diag}(x_1, \dots, x_N)$ and $Y = \text{diag}(y_1, \dots, y_N)$ with distinct increasing entries, normalized Haar measure on the unitary group satisfies

$$\int_{U(N)} \exp(\text{Tr}(XUYU^*)) dU = \frac{\prod_{r=1}^{N-1} r!}{\Delta(x)\Delta(y)} \det(e^{x_i y_j})_{i,j=1}^N, \quad (178)$$

where $\Delta(x) = \prod_{i < j} (x_j - x_i)$ and similarly for $\Delta(y)$ [29, 30, 31]. Hence

$$\log D_T^\Delta = \log \Delta(x) + \log \Delta(y) - \sum_{r=1}^{N_T-1} \log(r!) + \log \mathcal{I}_T, \quad (179)$$

with $\mathcal{I}_T$ the unitary integral in (178).

The large powers of T hidden separately in the two Vandermonde factors and the factorial product cancel at leading logarithmic scale. What remains is a genuine free-energy problem for two triangular empirical measures. The Baker separation theorem of Section 34 gives a polynomial lower bound for the spacings in $\Delta(x)$, so the row Vandermonde cannot collapse through anomalously close conditional exponents. What is not known is a bound on the full combination (179) strong enough to beat the primewise divisor.

Research target 37.6 (HCIZ–valuation comparison). Determine the T^4 -scale limit, or a sufficiently sharp upper bound, of

$$\log D_T^\Delta - \log Q_T^\Delta$$

for the conditional triangular row lattice and balanced mixed columns. It would suffice to prove that this quantity is negative for all sufficiently large T . Corollary 37.4 shows that such a result must use either the HCIZ/Vandermonde cancellation in (179), extra p -adic cancellation beyond the assignment minima, or both.

This formulation has a useful advantage over a generic interpolation determinant: both limiting measures and all lattice densities are explicit. It is therefore suitable for numerical free-energy reconnaissance before attempting a proof.

38 The first tied p -adic layer is lower order

Theorem 36.6 says that extra divisibility must come from ties among minimum-valuation Leibniz terms. Ties are especially abundant in the cross-coprime branch (159): at $p = 2$ all columns with the same u have the same valuation profile, and at $p = 3$ all columns with the same v do. It is natural to hope that cancellation inside those large fibers supplies the missing leading-order fraction.

The present section computes that first cancellation layer exactly. It factors into ordinary Vandermondes of the odd row parts. There are two useful bounds. Standard upper bounds for p -adic linear forms in two logarithms give the sharper $O(T^{3/2}(\log T)^C)$ estimate. More importantly for formalization, the elementary fact $p^{v_p(z)} \leq |z|$ already gives an unconditional $O(T^2)$ bound. Either is lower order than the mixed determinant height $\asymp T^{5/2}$. This does not bound the final determinant valuation—higher tropical layers may cancel again—but it rules out a one-step fiber explanation of the missing fraction without requiring deep p -adic approximation theory.

38.1 Tropical initial form at two

Assume throughout this section that $\gcd(MA, 6) = 1$. For the block rows, write

$$q_k = a_{T,k} = 3^{n_k} A^k, \quad n_k = T - \lfloor k\beta \rfloor. \quad (180)$$

Both M and every q_k are odd, and

$$m_{T,k}^u a_{T,k}^v = 2^{un_k} M^{ku} q_k^v. \quad (181)$$

For each integer $u \geq 0$, let

$$V_u = \{v : (u, v) \text{ is one of the selected } N \text{ columns}\}, \quad s_u = \#V_u.$$

Because columns are selected by the inequality $u + \vartheta v \leq L_N$ up to one boundary shell, each nonempty V_u is an initial interval of integers. Sort the distinct u values in increasing order. Since $n_0 > n_1 > \dots > n_{N-1}$, every minimizing assignment for the cost un_k sends one consecutive row block R_u of size s_u to the columns in the fiber $\{u\} \times V_u$.

Let $\tau_{2,T}$ be the minimum from (142). The sum of all Leibniz terms having exactly this valuation is the 2-adic tropical initial form. Up to an overall sign and an odd factor, it is

$$I_{2,T} = \prod_u \det(q_k^v)_{k \in R_u, v \in V_u} = \prod_u \prod_{\substack{i < j \\ i, j \in R_u}} (q_j - q_i). \quad (182)$$

The first equality follows by grouping the optimal permutations by fibers and factoring the odd number M^{ku} from each row in a fixed u block; the second is the ordinary Vandermonde identity because V_u is consecutive.

Proposition 38.1 (Exact first 2-adic initial form). *There is an integer $Z_{2,T}$ such that*

$$2^{-\tau_{2,T}} D_T = \varepsilon_{2,T} I_{2,T} + 2^{\lfloor \beta \rfloor} Z_{2,T}, \quad (183)$$

where $\varepsilon_{2,T}$ is odd. Consequently

$$v_2(D_T) - \tau_{2,T} \geq \min\{v_2(I_{2,T}), \lfloor \beta \rfloor\}. \quad (184)$$

No equality assertion follows without controlling the valuation of the full outside-fiber remainder $Z_{2,T}$.

Proof. The rearrangement inequality with repeated column weights identifies the minimizing permutations exactly as above, and summing their signed unit parts gives (182). Any nonminimizing assignment must move at least one column across two adjacent u fibers. An exchange across the corresponding boundary changes the cost by

$$(u' - u)(n_i - n_j).$$

Here $u' - u \geq 1$, while consecutive row depths differ by $\lfloor (k+1)\beta \rfloor - \lfloor k\beta \rfloor \geq \lfloor \beta \rfloor$. Hence every nonminimal Leibniz term has valuation at least $\tau_{2,T} + \lfloor \beta \rfloor$, proving the expansion. The valuation statement follows from the ultrametric inequality. $\square$

38.2 The symmetric initial form at three

Put

$$r_k = m_{T,k} = 2^{n_k} M^k. \quad (185)$$

At $p = 3$ the entry factors as

$$m_{T,k}^u a_{T,k}^v = 3^{vn_k} r_k^u A^{kv}.$$

For a selected value of v , let U_v be the consecutive set of selected u values and let R'_v be the corresponding consecutive block of rows in the minimum assignment. The same argument gives

$$I_{3,T} = \prod_v \prod_{\substack{i < j \\ i, j \in R'_v}} (r_j - r_i) \quad (186)$$

and

$$3^{-\tau_{3,T}} D_T = \varepsilon_{3,T} I_{3,T} + 3^{\lfloor \beta \rfloor} Z_{3,T}, \quad (187)$$

with $\varepsilon_{3,T}$ a 3-adic unit.

38.3 Two-logarithm bounds inside the fibers

For $i < j$, set $d = j - i$ and $h = n_i - n_j = \lfloor j\beta \rfloor - \lfloor i\beta \rfloor$. The common factors in the row values are units at the relevant place, so

$$v_2(q_j - q_i) = v_2(A^d - 3^h), \quad (188)$$

$$v_3(r_j - r_i) = v_3(M^d - 2^h). \quad (189)$$

Neither difference vanishes: $A^d = 3^h$ or $M^d = 2^h$ would give $d\beta = h$, contrary to the irrationality of β .

A standard consequence of p -adic linear-form theory is that for fixed multiplicatively independent positive rationals α_1, α_2 and a fixed prime p , there are effectively computable constants C_0, C_1 such that

$$v_p(\alpha_1^d - \alpha_2^h) \leq C_0 + C_1(\log(2 + \max\{d, h\}))^2. \quad (190)$$

Much sharper versions, with linear dependence on $\log B$, are known in important two-logarithm regimes; the square is more than sufficient here [21, 22, 23]. Apply this to $(A, 3, p = 2)$ and $(M, 2, p = 3)$. Since $d, h = O_\beta(N)$, each factor in (182)–(186) has valuation $O_{M,A}((\log N)^2)$.

For the qualitative scale separation, however, no linear-form theorem is needed. If p is prime and $a^d \neq b^h$, the divisor-size inequality gives the explicit finite bound

$$v_p(a^d - b^h) \leq 1 + d \lceil \log_p a \rceil + h \lceil \log_p b \rceil. \quad (191)$$

Indeed, $p^{v_p(a^d - b^h)}$ divides the nonzero difference, while $|a^d - b^h| \leq a^d + b^h$ and each power is bounded by the corresponding power of p . Inside a Beatty row block, $h \leq B_\beta d$ for a fixed integer B_β , so the right side is $O_{M,A}(d)$. A block of size s then contributes $O_{M,A}(s^3)$, because $\sum_{0 \leq i < j < s} (j - i) = O(s^3)$. There are $O(L_N)$ fibers of size $O(L_N)$, hence the total elementary budget is

$$O_{M,A}(L_N^4) = O_{M,A}(N^2) = O_{M,A}(T^2). \quad (192)$$

The number of pairs within the vertical fibers is

$$\sum_u \binom{s_u}{2} = \frac{L_N^3}{6\vartheta^2} + O_\vartheta(L_N^2) = O_\vartheta(N^{3/2}), \quad (193)$$

while the number within the horizontal fibers is

$$\sum_v \binom{\#U_v}{2} = \frac{L_N^3}{6\vartheta} + O_\vartheta(L_N^2) = O_\vartheta(N^{3/2}). \quad (194)$$

We obtain the following rigorous scale separation.

Theorem 38.2 (First tied layer is lower order). *In the cross-coprime branch,*

$$v_2(I_{2,T}) + v_3(I_{3,T}) = O_{M,A}(N^2) = O_{M,A}(T^2). \quad (195)$$

If one invokes the quoted two-logarithm estimate (190), the stronger bound $O_{M,A}(N^{3/2}(\log N)^2)$ also holds. In particular,

$$\frac{v_2(I_{2,T}) \log 2 + v_3(I_{3,T}) \log 3}{H_T} \longrightarrow 0. \quad (196)$$

Proof. For the elementary assertion, sum (191) over the pairs in each fiber and use (192). For the sharper assertion, sum (190) over the pair counts (193) and (194), using (188)–(189). Finally use $N \asymp_\beta T$ and $H_T \asymp_\beta T^{5/2}$. $\square$

Status: [\[kernel-verified Lean\]](#) for the exact finite valuation and cube-budget skeleton; [\[classical/open background\]](#) only for the optional sharper two-logarithm estimate.

The module `BlockVandermondeValuationBudget` proves the exact collision-depth identity `padicValNat_natVandermondeProduct_eq_sum_collisionCount`, the elementary power-difference bound `padicValNat_natAbs_pow_sub_pow_le`, and the variable-block and power-difference cube bounds `padicValNat_blockVandermondeNatProduct_le_cubeBudget` and `padicValNat_blockVandermondeNatProduct_le_powDifferenceCubeBudget`. The displayed $O(T^2)$ specialization still uses the elementary paper-level count of Beatty fiber sizes; the finite inequality on which it rests is kernel-verified. The sharper $O(T^{3/2}(\log T)^2)$ statement uses the cited classical estimate and is not claimed as a Lean theorem.

38.4 What a leading excess would have to do

Theorem 38.2 does not prove $v_2(D_T) - \tau_{2,T} = o(H_T)$ or its 3-adic analogue. If the initial form is divisible by a power exceeding the first tropical gap, terms from the next valuation layer enter; those can cancel, exposing another layer, and so on. A leading-order excess is therefore possible only through control of the valuations and cancellations in the truncated sums at cutoffs growing with T . It could arise from unusually divisible low-excess minor products, cancellation within one layer, or carrying across several layers. Cancellation confined to the minimum-permutation fibers is quantitatively negligible. Formula (82) makes this statement finite and exact: at depth K one need only control the signed products of unit minors attached to row partitions of excess less than K . The unit-adapted bases above extract a universal divisor from the relevant univariate minors, but its paper-level blockwise scale is only $O(N^{3/2})$. What is missing is an upper bound on the valuation of the residual truncated sum, uniform for cutoffs growing with T .

This yields a more precise version of the p -adic research target.

Research target 38.3 (No shallow-layer proof). Either prove that the tropical cancellation cascade has total depth $o(H_T)$ —which would retire the single-block p -adic strategy—or identify a structural identity forcing $\gg H_T$ cumulative cancellation across successively nonminimal assignments. Pairwise congruences such as (188) and (189) cannot by themselves supply the latter, because p -adic logarithm bounds make each congruence only polylogarithmic in the row separation.

38.5 Exact finite cascade reconnaissance

Exact integer experiments help rule out two tempting but false simplifications of Target 38.3. For synthetic cross-coprime parameters M, A , put

$$\beta = \log_2 M, \quad n_k = T - \lfloor k\beta \rfloor,$$

and take the first $N = 1 + \lfloor (T+1)/\beta \rfloor$ exponent pairs (u, v) in increasing order of $u + \vartheta v$. The tested determinant is exactly

$$D_T = \det \left((2^{n_k} M^k)^u (3^{n_k} A^k)^v \right)_{0 \leq k < N, (u,v)}.$$

Let $e_p = v_p(D_T) - \tau_p$, let I_p denote the first-fiber Vandermonde product, and let U_p be the common row-block ordering exponent determined by the column-fiber sizes. The weights come from `AllLayerUnitBasis`, their finite sums from `UnitOrderingGain`, and their uniform contribution to every row-block term and the full determinant from `RowBlockUnitDivisibility`. For $(M, A) = (5, 7)$, exact DVR elimination gives:

T	N	e_2	U_2	$v_2(I_2)$	e_3	U_3	$v_3(I_3)$	$(e_2 \log 2 + e_3 \log 3)/H_T$
30	14	26	17	29	14	9	14	0.0360124
100	44	157	136	157	87	74	87	0.0101461
300	130	868	796	870	517	471	530	0.00353027

The data are compatible with an $N^{3/2}$ excess and incompatible with no proved asymptotic; their purpose is diagnostic.

Two exact small cases are more informative. At $(M, A, T, p) = (5, 7, 30, 2)$ one has

$$\tau_2 = 177, \quad v_2(D_T) = 203, \quad e_2 = 26 < 29 = v_2(I_2).$$

Writing $2^{-\tau_2} D_T = \sum_{\delta} 2^{\delta} C_{\delta}$, the first layer has $v_2(C_0) = 29$, whereas the layer $\delta = 6$ has $v_2(C_6) = 20$ and therefore smaller effective valuation 26. A nonminimal layer can genuinely beat the initial form. Conversely, at $(M, A, T, p) = (11, 5, 20, 2)$ the layers $\delta = 0$ and $\delta = 4$ both have effective valuation 9, but their sum is divisible by 2^{10} ; all other layers have effective valuation at least 12, and the final excess is 10. Thus even “take the minimum effective layer” is false because different layers can carry into one another.

Status: [\[exact computation\]](#) — all displayed determinant, assignment, layer, and valuation integers use exact integer arithmetic; the logarithmic ratios are ordinary floating approximations used only diagnostically. The no-dependency scripts `scripts/cascade_experiments.py`, `scripts/cascade_layer_report.py`, and `scripts/cascade_verify.py` reproduce the computations; the last directly Bareiss-checks the same twenty small determinants and agrees on their valuations.

The parameters are deliberately synthetic: they satisfy $\gcd(6, MA) = 1$, but do not assert the unknown identity $A = 3^{\log_2 M}$. The decreasing ratios and the apparent $N^{3/2}$ scale are evidence only. The two finite counterexamples to naive layer formulas are exact for the displayed mixed determinant family.

39 Synthesis of the determinant audits

The preceding calculations put several superficially different determinant proposals into the same currency: the fraction of the leading archimedean budget supplied automatically by valuations of the entries. The numerical constants should not be compared as though the underlying determinants were identical—the ordinary, mixed-block, and complete-grid systems have different sizes and different upper bounds—but they do show which mechanism is closest to a contradiction before any genuine cancellation is used.

Three conclusions survive every normalization.

Table 7: Leading termwise-divisor ceilings obtained in this continuation.

System	Ceiling or exact limit	Meaning
Ordinary input/output Vandermonde product	$1 - \frac{\delta_2 \log 2 + \delta_3 \log 3}{3\beta \log 6} < 0.8710491$	Exact fixed-support divisor fraction; all primes already occurring in M and A included.
Mixed single block, arbitrary branch	$1 - \frac{3}{8} \sum_p \mathfrak{G}(c_p) \leq 4/5$	Exact all-prime tropical limit and universal primitive-depth ceiling.
Mixed single block, cross-coprime	$\leq 7/10$	Uses $c_2 = (-1, 0)$, $c_3 = (0, -1)$ and the external sum $(1, 1)$.
Mixed single block, only structural primes in the cross-coprime branch	$3/10$	Exact contribution of $p = 2, 3$ before external primes or cancellation.
Complete triangular row grid	$\leq \pi/4 = 0.7853982\dots$	All-prime common divisor versus the sharp largest-Leibniz-term transport upper bound.
First tied fibers at 2 and 3	$o(1)$ of the mixed-block budget	Verified finite budgets plus elementary paper-level Beatty counting give $O(T^2)$ against $H_T \asymp T^{5/2}$; quoted two-logarithm estimates sharpen this to $O(T^{3/2}(\log T)^2)$.

1. Merely adding every prime in the fixed support does not close any determinant. It does, however, improve the diagnosis substantially: the universal mixed-block ceiling is $4/5$, not a heuristic “about one half”.
2. The exact primitive theorem is quantitatively useful. The condition $\min(v_2(M), v_3(A)) = 0$ creates one structural slope vector of norm at least $4/15$; the conservation law forces a second norm contribution of at least the same size.
3. A successful p -adic proof must be a cancellation theorem, not a divisibility-by-every- term theorem. In the easiest branch, even the first large tied family contributes only a lower-order amount. The required phenomenon would have to repeat through many tropical layers or interact with a much sharper archimedean free energy.

39.1 Correction of the post-attachment “one half / two thirds” note

After the attached unified report was produced, the repository acquired a short exploratory note (merged into Section 21 of this report) at commit `ff9a9f5aa3a3f02599de93fc053e60918f269855`. That note correctly identified the distinction between termwise min-plus divisibility and cancellation beyond the tropical minimum, and its real/ p -adic compatibility diagnosis remains valuable. It proposed, however, a structural $1/2$ ceiling and an all-prime $2/3$ ceiling. Neither constant is valid in the stated generality.

The structural averaging argument omitted the base and cross depths that appear in (150)–(151). Even under $v_2(M) = 0$, the terms $v_2(A)kv$ and $v_3(M)ku + v_3(A)kv$ do not disappear. The exact structural contribution is obtained by retaining c_2 and c_3 in Theorem 36.3; it can exceed one half of

H_T .

The all-prime step used the assertion that an oppositely sorted nonnegative triangular weight has assignment minimum at most two thirds of its mean. For the weight $Y = X + Z$ on the standard triangle, the exact values from Lemma 36.2 are

$$\mathbb{E}Y = \frac{2}{3}, \quad \Phi(1, 1) = \frac{4}{15}, \quad \frac{\Phi(1, 1)}{\mathbb{E}Y/2} = \frac{4}{5}.$$

Thus $4/5$, not $2/3$, is already attained by the relaxed external rank-one grading. The Gini-norm identity explains the general case and leads to the rigorous ceiling in Theorem 36.6. This is a correction of a calculation in an exploratory note, not a criticism of the source report itself; the attached unified report did not contain the claim.

Status: **[paper proof in this report]** — the correction follows from the exact transport formula, not from numerical experimentation.

40 Directed-rounding extension through 2^{27}

The attached report supplied a complete MPFR scan through 2^{26} . Because a finite extension is independent of the unresolved transcendence barrier and gives a useful stronger floor for every hypothetical generator, I extracted the literal C listing from its Appendix “Directed-rounding verifier”, compiled it against MPFR 4.2.2, replayed an old audit chunk, and scanned the next full dyadic range.

40.1 Manifest audit

The literal extracted source has SHA-256 digest

`f888c0d89be9366f01e2d26f339dcbe446556d39d4c5dbcf8f579ec0d0052e1d.`

This does *not* equal the digest

`ec709623f812478fbb7fb5567aa59fbfe3df0ae639ca366bc7c72e278bcfee2f`

printed in the source of that section. None of the ordinary line-ending, final-newline, or tab normalizations explains the discrepancy. The report therefore contains a stale or externally generated source-identity hash.

The semantic reproducibility check succeeds. Replaying the old interval $[29,360,128,33,554,432)$ at 192 bits gives exactly the report’s records

lower: $1.4439124872349507 \times 10^{-19}$ at $(30,103,832,713,403,245,915)$,
upper: $5.4132120119800561 \times 10^{-21}$ at $(29,411,704,687,581,893,443)$,

with zero failures. Thus the listing used here reproduces the old executable’s mathematical behavior on an audited four-million-input overlap even though the printed file digest does not identify it.

Status: [\[software computation\]](#) — this section is a directed-rounding software result. It is stronger numerically than the Lean theorem but is not a kernel proof. The manifest hash mismatch is explicitly part of the trust boundary.

40.2 The new range

The half-open interval $[2^{26}, 2^{27})$ was divided into sixteen chunks of width 2^{22} . The executable used 192-bit MPFR endpoints and outward rounding exactly as in the source report. There are 2^{26} integers in the range; the first, 2^{26} , is the sole power of two and is intentionally skipped. Hence the required nonpower count is $2^{26} - 1 = 67,108,863$.

The concatenated canonical extension log has SHA-256 digest

b11843448cd8946442aba13758b5b666e4ac6eae50d2243fbdc32d83274cea8d.

Every chunk reports zero failures, and the checked counts sum to the required total.

Table 8: Directed-rounding extension on $[2^{26}, 2^{27})$.

Chunk	Integer range	Checked	Smallest lower gap	Smallest upper gap
0	67,108,864–71,303,167	4,194,303	4.2837×10^{-20}	1.7723×10^{-19}
1	71,303,168–75,497,471	4,194,304	2.2961×10^{-20}	1.1433×10^{-20}
2	75,497,472–79,691,775	4,194,304	1.2813×10^{-19}	1.6063×10^{-20}
3	79,691,776–83,886,079	4,194,304	2.5420×10^{-20}	5.2012×10^{-20}
4	83,886,080–88,080,383	4,194,304	2.6166×10^{-20}	3.8496×10^{-20}
5	88,080,384–92,274,687	4,194,304	1.0923×10^{-19}	5.7144×10^{-20}
6	92,274,688–96,468,991	4,194,304	1.1017×10^{-20}	2.3227×10^{-20}
7	96,468,992–100,663,295	4,194,304	3.0759×10^{-20}	7.4810×10^{-20}
8	100,663,296–104,857,599	4,194,304	5.6465×10^{-20}	6.4744×10^{-20}
9	104,857,600–109,051,903	4,194,304	8.8144×10^{-20}	8.3662×10^{-21}
10	109,051,904–113,246,207	4,194,304	3.8892×10^{-20}	7.0327×10^{-21}
11	113,246,208–117,440,511	4,194,304	9.9365×10^{-21}	2.1814×10^{-20}
12	117,440,512–121,634,815	4,194,304	4.4964×10^{-20}	5.4132×10^{-21}
13	121,634,816–125,829,119	4,194,304	1.0774×10^{-20}	3.4093×10^{-20}
14	125,829,120–130,023,423	4,194,304	4.3609×10^{-20}	8.3073×10^{-22}
15	130,023,424–134,217,727	4,194,304	5.9082×10^{-21}	5.1326×10^{-20}
total		67,108,863		

40.3 New extremal records and high-precision replay

The smallest lower logarithmic margin in the extension is

$$5.9082259992879571 \times 10^{-21} \quad \text{at} \quad (m, n) = (132,833,765, 7,501,348,219,569). \quad (197)$$

A 120-decimal-digit reevaluation gives

$$m^\vartheta - n = 6.3939754533470409837680125510 \dots \times 10^{-8}.$$

The smallest upper logarithmic margin is the new global record

$$8.3073231820333117 \times 10^{-22} \quad \text{at} \quad (m, n) = (129, 799, 345, 7, 231, 571, 292, 851), \quad (198)$$

where

$$n + 1 - m^\vartheta = 8.6669904355818956853297952288 \dots \times 10^{-9}.$$

Both one-input cases were rerun independently at 320 and 512 bits. The directed margins and nearest integer were unchanged in every displayed digit, and both replays reported zero failures.

Theorem 40.1 (Software-level finite exclusion through 2^{27}). *At the stated MPFR software-trust level, for every integer $1 \leq m < 2^{27}$,*

$$m^\vartheta \in \mathbb{N} \implies m \text{ is a power of two.}$$

Consequently every nonintegral real x satisfying $2^x, 3^x \in \mathbb{Z}$ obeys

$$x > 27. \quad (199)$$

Proof. The attached report’s complete scan covers $1 \leq m < 2^{26}$, and the sixteen new chunks cover $2^{26} \leq m < 2^{27}$. In each case the directed logarithmic interval for $\vartheta \log m$ lies strictly between the intervals for the logarithms of two consecutive integers, except when m is a power of two. If $m = 2^x$ came from a nonintegral solution and $m < 2^{27}$, then $m^\vartheta = 3^x$ would be an integer at a non-power input, contradicting the scan. The endpoint $m = 2^{27}$ corresponds to the integral exponent 27, so a nonintegral solution must have $m > 2^{27}$ and hence $x > 27$. $\square$

This raises the software-level spacing ratio in the optimized finite-difference comparison from $26 \log 3$ to $27 \log 3 = 29.662531 \dots$. It does not change any kernel-verified theorem; the formal floor remains $x \geq 12$ until the finite certificates are replayed in Lean.

41 Additional proof routes pushed to their failure points

The determinant calculations were the main new line of attack, but they were not the only one. This section records four other routes in enough detail to prevent the same hidden loss from being rediscovered under different notation.

41.1 Transferring convergents of β to ϑ

Let p/q be a continued-fraction convergent of a hypothetical primitive generator β and put

$$\varepsilon = q\beta - p.$$

For all sufficiently large q , $|\varepsilon| < 1/2$. Define the two nonzero integer gaps

$$D_2 = M^q - 2^p = 2^p(2^\varepsilon - 1), \quad D_3 = A^q - 3^p = 3^p(3^\varepsilon - 1). \quad (200)$$

They manufacture an exact rational approximation to ϑ :

$$R_{p,q} := \frac{2^p D_3}{3^p D_2} = \frac{3^\varepsilon - 1}{2^\varepsilon - 1} \in \mathbb{Q}. \quad (201)$$

Taylor expansion at zero gives

$$R_{p,q} = \vartheta + \frac{\log 3 (\log 3 - \log 2)}{2 \log 2} \varepsilon + O(\varepsilon^2), \quad (202)$$

so $|R_{p,q} - \vartheta| \ll |\varepsilon|$.

This looks promising because ϑ has an explicit irrationality measure. Wu’s simultaneous linear-independence estimate for $1, \log 2, \log 3, \log 5$ gives, after setting the unused coefficients to zero, a conservative exponent

$$\left| \vartheta - \frac{r}{s} \right| \geq c_{\vartheta} s^{-\mu_{\vartheta}}, \quad \mu_{\vartheta} = 8.616 \quad (203)$$

for all sufficiently large denominators s [24]. The obstacle is the denominator created by (201). Before reduction it is at most

$$3^p |D_2| \ll 3^p 2^p |\varepsilon| = 6^p |\varepsilon|. \quad (204)$$

Combining (202)–(204) with (203) gives only

$$|\varepsilon| \gg 6^{-p\mu_{\vartheta}/(1+\mu_{\vartheta})}. \quad (205)$$

Since

$$6^{-\mu_{\vartheta}/(1+\mu_{\vartheta})} = 2^{-2.316\dots},$$

this is substantially weaker than the elementary bound $|\varepsilon| \gg 2^{-p}$ obtained directly from $D_2 \neq 0$. Sharper real irrationality measures cannot repair the architecture: every irrationality exponent is at least two, and the unreduced denominator already contains both 2^p and 3^p .

Proposition 41.1 (Exact denominator-loss diagnosis). *The convergent-transfer construction can improve the elementary gap only if one proves an exponential lower bound for*

$$G_{p,q} = \gcd(2^p D_3, 3^p D_2) \quad (206)$$

that survives reduction of $R_{p,q}$. In the absence of such a bound, every known irrationality measure for ϑ is defeated by the factor 6^p in (204).

No mechanism forcing $G_{p,q}$ to be exponentially large was found. The two gaps share the same small real parameter ε , but real closeness does not impose a common prime divisor.

41.2 Two-row determinants and fixed-support S -units

For two block rows $i < j$, put $d = j - i$ and $h = n_i - n_j = \lfloor j\beta \rfloor - \lfloor i\beta \rfloor$. Their 2×2 determinant has the exact factorization

$$m_{T,i} a_{T,j} - m_{T,j} a_{T,i} = 6^{n_j} (MA)^i (2^h A^d - 3^h M^d). \quad (207)$$

Moreover

$$\frac{2^h A^d}{3^h M^d} = \left(\frac{3}{2} \right)^{\beta d - h}. \quad (208)$$

Thus a good rational approximation h/d to β makes the bracket in (207) a small *relative* difference of two integers supported on the fixed prime set dividing $6MA$.

Let

$$U = 2^h A^d, \quad V = 3^h M^d, \quad G = \gcd(U, V).$$

The reduced nonzero integer $(U - V)/G$ gives

$$|\beta d - h| \gg_{M,A} \frac{G}{\max\{U, V\}}. \quad (209)$$

The exponent of the right-hand side can be written explicitly as the positive part of a valuation-vector difference:

$$\log \frac{U}{G} = \sum_p \max\{h \mathbf{1}_{p=2} + dv_p(A) - h \mathbf{1}_{p=3} - dv_p(M), 0\} \log p. \quad (210)$$

As $h/d \rightarrow \beta$, this is $\kappa_{M,A}d + o(d)$ for a positive constant $\kappa_{M,A}$. Hence (209) is merely exponential in d , much weaker than the Baker–Matveev polynomial separation of Section 34.

The *abc* conjecture applied to the reduced equation $U/G - V/G = (U - V)/G$ would make the difference nearly as large as a fixed positive power of U/G , but it would still allow a relative gap $e^{-\epsilon d}$. Continued-fraction gaps of order $1/d$ are much larger. Thus this direct *S*-unit equation does not become a proof even after inserting *abc*; one would need an additional theorem controlling the prime support or radical of the difference.

41.3 Automaticity and the synchronized Sturmian word

The coding $s_k = \lfloor (k+1)\alpha \rfloor - \lfloor k\alpha \rfloor$ from Section 33 is aperiodic Sturmian. A standard theorem in automatic-sequence theory says that no aperiodic Sturmian word is automatic in any integer base [27, 28]. One quick obstruction is frequency: the frequency of 1 in s is the irrational number α , whereas an automatic sequence having a letter frequency has rational frequency.

This observation blocks a tempting but invalid shortcut. The rationality of $U = 2^\alpha$ does *not* make the threshold itinerary of $R \mapsto UR/2^{\mathbf{1}_{UR \geq 2}}$ finite-state; its exact rational denominators grow with time. The same is true for V in base three. Therefore Cobham’s theorem cannot be invoked merely because the same word is produced by a dyadic and a triadic rational system. A viable automata proof would first have to manufacture genuine finite kernels in two multiplicatively independent bases, and doing so is already a new arithmetic theorem about U and V .

41.4 Real and p -adic logarithms

At the real place the counterexample is exactly

$$\log M \log 3 - \log A \log 2 = 0. \quad (211)$$

At $p = 2$ or 3 , after stripping valuation and torsion parts, the units among $M, A, 2, 3$ admit p -adic logarithms. The determinant entries then have unit parts whose p -adic logarithms are bilinear forms in the row and column coordinates. This is the natural language for the cancellation cascade isolated in Section 38.

What is missing is a transfer principle from (211) to a congruence or rank drop among those p -adic logarithms. Real exponentiation by the transcendental number β does not define a compatible p -adic exponent, and the equality $M = 2^\beta$ has no p -adic meaning. Known p -adic linear-form theorems control a nonzero form once its integer coefficients and algebraic bases are specified; they do not create a p -adic form from a real one.

Research target 41.2 (Real/ p -adic compatibility). Find an algebraic construction, valid for the integer pair (M, A) alone, that converts (211) into either

1. a nontrivial congruence among p -adic unit logarithms with precision $\gg T$, repeated over $\gg T^{3/2}$ fiber pairs; or
2. a rank drop in the tropical initial matrices persisting through $\gg T$ valuation layers.

Without such a construction, ordinary p -adic logarithm bounds point in the opposite direction: they show that individual congruences are only polylogarithmically deep.

41.5 Transcendental-curve point counting

The graph $y = x^\vartheta$ is definable in the real exponential field, so Pila–Wilkie type results [32] give subpolynomial bounds for rational points outside algebraic pieces. Those generic bounds are far too weak here. The curve already contains the infinite sequence $(2^n, 3^n)$, giving $\asymp \log H$ integer points of height at most H , while a hypothetical second generator produces $\asymp (\log H)^2$ points globally. A proof by point counting would therefore need a sharp $O(\log H)$ theorem with a structural classification of the extremal sequence, not merely $O(H^\epsilon)$. Such a theorem would essentially be another form of the rank-one conclusion sought by the conjecture.

42 Exact sufficient conditions extracted from the new analysis

A useful continuation should not stop at saying that the present determinant fails. The calculations above isolate numerical thresholds which, if crossed by a future theorem, would settle the conjecture. This section records those thresholds as precise implications.

42.1 The mixed-block cancellation-excess criterion

Let

$$\mathcal{P} = \{p : p \mid 6MA\}$$

be the fixed support of all determinant entries, and retain the mixed block determinant D_T , its tropical minima $\tau_{p,T}$, and the archimedean scale H_T . Define the support-prime cancellation excess

$$\mathcal{E}_T = \sum_{p \in \mathcal{P}} (v_p(D_T) - \tau_{p,T}) \log p \geq 0 \quad (212)$$

and the exact tropical deficit

$$\delta(M, A) = \frac{3}{8} \sum_{p \in \mathcal{P}} \mathfrak{G}(c_p). \quad (213)$$

Theorem 36.3 says $\log Q_T^{\text{trop}} = (1 - \delta(M, A) + o(1))H_T$, while the elementary Leibniz upper bound is $\log D_T \leq H_T + o(H_T)$.

Theorem 42.1 (Cancellation-excess criterion). *If a hypothetical primitive counterexample satisfies*

$$\liminf_{T \rightarrow \infty} \frac{\mathcal{E}_T}{H_T} > \delta(M, A), \quad (214)$$

then no such counterexample exists. In particular it is enough to prove an excess greater than $H_T/5$ uniformly in all primitive branches; in the cross-coprime branch the exact threshold is at least $3H_T/10$.

Proof. The selected support primes alone give

$$\log D_T \geq \sum_{p \in \mathcal{P}} v_p(D_T) \log p = \log Q_T^{\text{trop}} + \mathcal{E}_T.$$

Under (214), the right-hand side exceeds H_T by a fixed positive multiple for all large T , contradicting the Leibniz upper bound. The numerical thresholds follow from Theorem 36.6 and Corollary 36.7. $\square$

This criterion is not tautological: it involves only valuations at the fixed finite set of primes already present in $6MA$, whereas D_T can acquire arbitrary new prime factors. It states exactly how much same-support cancellation is required.

42.2 A denominator-cancellation criterion

For the transferred approximant $R_{p,q}$ of Section 41.1, let $s_{p,q}$ be its reduced denominator. If one could prove along infinitely many convergents that

$$s_{p,q} \ll q^\eta \quad \text{for some } \eta < \frac{1 + \mu_\vartheta}{\mu_\vartheta}, \quad (215)$$

then Wu’s measure would give $|R_{p,q} - \vartheta| \gg q^{-\eta\mu_\vartheta}$, whereas (202) and the convergent bound give $|R_{p,q} - \vartheta| \ll q^{-1}$. Since $\eta\mu_\vartheta < 1 + \mu_\vartheta$ is not by itself enough, one tracks $|\varepsilon| < 1/q_+$ more precisely; a contradiction follows whenever $q_+ \gg q^{\eta\mu_\vartheta + \epsilon}$ along the same subsequence. A simpler sufficient condition is

$$s_{p,q} = q^{o(1)} \quad (216)$$

for infinitely many convergents with $q_+ \geq q^{1+\epsilon}$.

Equivalently, the gcd $G_{p,q}$ in (206) would have to cancel essentially the entire exponential factor $6^p|\varepsilon|$, leaving a subpolynomial denominator. This is a sharp formulation of the otherwise vague request for “correlation between the two integer gaps.”

42.3 An automaticity criterion

The synchronized coding gives another exact, if presently inaccessible, implication.

Proposition 42.2 (Finite-kernel criterion). *If under a hypothetical counterexample the common word $s_k = \lfloor (k+1)\alpha \rfloor - \lfloor k\alpha \rfloor$ were automatic in any integer base $q \geq 2$, then the Alaoglu–Erdős conjecture would follow.*

Proof. An automatic Sturmian word is ultimately periodic. Its slope, equal to the frequency of the letter 1, is therefore rational. But $\alpha = \beta - \lfloor \beta \rfloor$ would then be rational, and the rational-exponent theorem would make β integral. $\square$

The missing step is not symbolic dynamics; it is proving automaticity from the two rational multipliers U and V . The growing denominators in (104) show exactly why the obvious state space is infinite.

43 Branch-sensitive research program

The source report already classifies a hypothetical pair much more sharply than “two integers on a curve.” The next attack should exploit that classification rather than seek a single estimate that treats every branch identically.

43.1 Branch I: independent external valuation directions

Here the second-iterate kernel K_β is zero, equivalently the external prime-valuation vectors of M and A are independent. Corollary 36.5 gives a strict angular defect $\Delta_{\text{ext}} > 0$ in the all-prime tropical fraction. That makes the termwise divisor farther from the archimedean upper bound, but the same independence also supplies at least two genuinely different external primes at which cancellation can be studied.

The concrete target is a simultaneous p -adic rank theorem for the tropical initial matrices: show that if cancellation persists through L layers at two non-collinear external primes, then an integer relation appears among the two external valuation vectors. Such a theorem would contradict $K_\beta = 0$. Existing one-prime linear-form bounds do not compare the residue kernels at different primes; this is the genuinely new ingredient required.

43.2 Branch II: a common external valuation ray

When $K_\beta \neq 0$, all external vectors are collinear and the source report gives the exact Möbius-orbit condition

$$\beta = \frac{u + v\vartheta}{r + s\vartheta} \quad (u, v, r, s \in \mathbb{Q}, us - vr \neq 0). \quad (217)$$

This is the branch in which the external triangle inequality can be an equality. It is also arithmetically lower-dimensional: after clearing denominators, the external parts of M and A are powers of one common primitive core.

The best next step is therefore not a generic all-prime determinant. It is a mixed real/radical invariant built from the canonical core and the two binomial fields already formalized in the repository. Useful intermediate targets are:

1. compare the field norms of the canonical radicals on the two sides of (217);
2. show that simultaneous radical degrees force an additional common perfect-power degree, contradicting affine primitivity; or
3. prove that the Möbius coefficients in (217) cannot have the sign pattern imposed by nonnegative prime valuations unless one output is $\{2, 3\}$ -smooth.

The repository modules

CanonicalRadicalFieldNorm, CanonicalRadicalValuation,
CanonicalRadicalSharpness, CanonicalRadicalDivisibleHull,
SimultaneousRadicalDegrees

provide much of the algebraic plumbing for this branch.

43.3 Branch III: one smooth output

If M is $\{2, 3\}$ -smooth, then $\beta = a + d\vartheta$ with $d > 0$; if A is smooth, there is a symmetric affine formula. The source corpus proves that both outputs cannot be smooth and that the other output contains a prime at least five. In the extreme case $M = 3^d$, the problem is the localized radical assertion

$$(3^d)^\vartheta \notin \mathbb{Z}[1/3].$$

This branch is the smallest transcendence-theoretic wall and deserves separate treatment. Potentially useful objects are the irreducible binomial for the least localized radical, its field norm, and reduction at primes over two and three. A successful local argument must remain sensitive to the real identity defining ϑ ; ordinary ideal factorization alone is already exhausted by the canonical-radical modules.

43.4 Branch IV: cross-coprime primitive outputs

The cross-coprime branch has the strongest determinant constants and the cleanest residue matrices. Here the exact first tropical forms are the fiber Vandermondes of Section 38. This is the best laboratory for determining whether the cancellation cascade is real or illusory.

The next computational experiment should not search for counterexamples. For synthetic odd pairs (M, A) and Beatty depths generated by $\beta = \log_2 M$, compute

$$v_2(D_T) - \tau_{2,T}, \quad v_3(D_T) - \tau_{3,T}$$

for increasing T , together with the histogram of tropical layers contributing at the final valuation. Small exact experiments performed during this work show an excess compatible with $N^{3/2}$ rather than $N^{5/2}$, matching Theorem 38.2; they are reconnaissance only, since a synthetic A does not satisfy $A = 3^\beta$. A proof should seek a uniform $o(H_T)$ bound, probably through p -adic exponential-polynomial zero estimates.

43.5 Archimedean priority: solve the triangular HCIZ limit

The complete grid finally supplies the missing T^2 rows, and its empirical measures are explicit. This makes Target 37.6 the highest-priority analytic problem. The required calculation is a large- N asymptotic for the HCIZ integral with both spectra having density $2t$ on $[0, 1]$, under the coupled scaling $x_i y_j \asymp T^2$ and $N \asymp T^2$. Even a nonsharp upper bound improving the largest-term energy by more than $1 - \pi/4$ would materially change the p -adic comparison.

A practical sequence is:

1. numerically evaluate $\log \det(e^{x_i y_j})$ using high-precision total-positive algorithms for moderate T ;
2. subtract the exact Vandermonde/factorial terms in (179) to estimate the free-energy limit;
3. derive the corresponding variational problem using known HCIZ large-deviation methods;
4. only then reinsert the primewise valuation lower bounds.

This route is attractive because the analytic limit does not depend on the unknown prime factorization of M and A ; only β enters the overall scaling.

43.6 Finite verification priority

The software floor is now $x > 27$. The immediate formal target is not another raw doubling of the MPFR scan but a replayable certificate architecture. A realistic hierarchy is:

1. a fast untrusted generator records only near-boundary pairs and rational interval data;
2. a small verified checker proves each interval enclosure and the monotone coverage between records;
3. chunks are combined by a theorem whose statement mentions only integer endpoints;
4. the final theorem is imported into the aggregate surface without `native\_decide` or an opaque external axiom.

The source-hash discrepancy found in Section 40 makes this especially worthwhile: mathematical replay is more robust than a manifest that identifies an external binary only by prose.

44 Lean formalization blueprint

The source corpus is already unusually strong on exact conditional structure. The most useful continuation is therefore not a second informal reimplementing of those results, but a small collection of sharply separated modules whose statements expose the genuinely new invariants. This section gives a dependency order designed to keep arithmetic, asymptotics, and transcendence inputs from leaking into one another.

44.1 Proposed module graph

Proposed module	Principal statement	Main dependencies
SturmianNormalization	Exact recurrences for R_k, S_k , common jump word s_k , and the exponent/output correspondence.	Existing primitive-generator and denominator-normalization modules; elementary floor arithmetic.
ReducedOrbitDenominators	Lowest-term denominators $q_k^{(2)}, q_k^{(3)}$ and the full-depth alternative from $\min(v_2(M), v_3(A)) = 0$.	SturmianNormalization ; unique factorization and valuations.
EffectiveGeneratorSeparation	An abstract interface turning a two-logarithm lower bound into $ \beta - p/q \geq cq^{-\mu}$ and a polynomial recurrence for continued-fraction denominators.	A stated external Baker–Matveev hypothesis, plus existing continued-fraction lemmas.
BlockVandermondeValuation	Exact v_2 and fixed-support v_p formulas for the conditional block Vandermonde; output analogue.	Existing exact block parametrization; valuation of a difference with unequal valuations.
OrdinaryAdelicFraction	The asymptotic fraction (131) and the universal $1 - \log 2/(3 \log 6)$ ceiling.	BlockVandermondeValuation ; polynomial floor sums; primitive-output depth theorem.
TriangleGini	$\mathfrak{G}$ is a norm and satisfies the explicit formula (147).	Elementary measure integration, or a finite polygonal integration library.

Proposed module	Principal statement	Main dependencies
FiniteTropicalAssignment	Exact rearrangement formula for a finite row list and a finite column-weight list.	Sorting/permutations; finite rearrangement inequality.
MixedColumnRiemannSums	Lattice count (133) and convergence of scaled low-weight columns to uniform measure on $\mathcal{T}$.	Two-dimensional lattice-point estimates and Riemann sums.
MixedAdelicGini	The limit (154), conservation $\sum_p c_p = 0$, and branch-sensitive three-vector bound.	Previous three modules; finite prime support.
MixedTropicalCeilings	Universal 4/5, cross-coprime 7/10, and structural 3/10 constants.	MixedAdelicGini ; primitive-output depth theorem; explicit values from TriangleGini .
FullGridTransport	Row/column empirical distributions, countermonotone energy $\pi/8$, maximum energy 1/2, and the ratio $\pi/4$.	Triangular lattice sums, quantile integration, rearrangement.
BeattyTropicalFiber [accepted]	Exact minimizing fibers, floor-sized gap, normalized quotient divisibility, and block-determinant expansion for Beatty rank-one matrices.	TropicalInitialForm , BlockFiberFactorization ; rearrangement and floor arithmetic.
RowBlockLaplaceExpansion [accepted]	Exact grouping of every assignment layer by row partitions; normalized and depth- K -truncated expansions with coefficients equal to signed products of unit minors.	BeattyTropicalFiber , BlockFiberFactorization ; finite fiberwise sums and determinant algebra.
AssignmentExcessGeometry [accepted]	Prefix-surplus crossing energy, the Beatty lower bound $c(a) - \tau \geq \lfloor \beta \rfloor \mathcal{E}(a)$, and for $N > 0$ the cutoff support bound $\#\{c(a) - \tau < K\} \leq \binom{N+q-1}{q}$.	RowBlockLaplaceExpansion ; finite rearrangement, summation by parts, and stars and bars.
BlockVandermondeValuationBudget [accepted]	Budget collision-energy identity, elementary power-difference bound, and variable-block cube budget; the finite input to the paper-level $O(T^2) = o(H_T)$ first-layer specialization.	BlockFiberFactorization ; elementary p -adic valuation and finite sums.
AllLayerUnitBasis [accepted]	Monic unittriangular and determinant-one column changes extract complete-determinant unit divisors on positive dyadic/triadic unit nodes, with p -ordering weights $j + v_2(j!)$ and $\lfloor j/2 \rfloor + v_3(\lfloor j/2 \rfloor!)$.	BlockVandermondeValuationBudget ; integer polynomial evaluation and Vandermonde determinants.
UnitOrderingGain [accepted]	Exact block-weight decompositions and finite quadratic/cubic aggregate bounds; at most S blocks of size at most S have dyadic gain at most S^3 and twice the triadic gain at most S^3 .	AllLayerUnitBasis ; Legendre digit sums and finite-sum bounds.
RowBlockUnitDivisibility [accepted]	Uniform dyadic/triadic unit divisors for every scaled row-block minor product and the additive full-determinant conclusion $p^{\tau+U} \mid \det$.	RowBlockLaplaceExpansion , UnitOrderingGain ; row-scaled Vandermonde reindexing and a determinant-one fiberwise basis change.
FiniteCheck2Pow27	A kernel-replayable certificate excluding non-powers of two below 2^{27} .	Existing finite-check certificate architecture; generated chunk certificates.

The first two modules are almost entirely combinatorial. The accepted Beatty, block-budget, and unit-basis modules likewise import no transcendence theory; in particular, the qualitative first-layer scale separation no longer needs an abstract p -adic logarithm interface. Any future use of the sharper estimate (190) should still be carried as a named classical input unless and until Mathlib contains a suitable theorem. This keeps the trust boundary explicit.

44.2 Finite statements before asymptotics

The continuum notation in Sections 36 and 37 is compact, but Lean should first see finite inequalities. For rows $0 \leq k < N$ and columns $z_j \in \mathbb{R}$, the finite rearrangement lemma should assert

$$\min_{\sigma \in S_N} \sum_{k=0}^{N-1} k z_{\sigma(k)} = \sum_{k=0}^{N-1} k z_{N-1-k}^{\uparrow}, \quad (218)$$

where $z^{\uparrow}$ is the nondecreasing rearrangement. Every tropical assignment in this report is an affine rescaling of (218). The limiting Gini identity can then be obtained by three auditable steps:

1. prove convergence of empirical distributions for the triangular lattice;
2. prove convergence of the first moments and pairwise absolute-difference moments;
3. invoke the finite identity

$$\frac{1}{N^2} \sum_{i,j} |z_i - z_j| = \frac{2}{N^2} \sum_{j=0}^{N-1} (2j - N + 1) z_j^{\uparrow}.$$

This route avoids formalizing optimal transport as a general theory.

44.3 Exact theorem sketches

The intended public statements can be phrased close to the mathematics.

```
/-- A primitive counterexample determines one common lower mechanical word
driving both normalized rational orbits. -/
theorem exists_synchronized_sturmian_data
  (hfail : ExistsNonintegralSolution) :
  Exists fun B : Nat => Exists fun alpha U V : Real =>
    0 < alpha /\ alpha < 1 /\ Irrational alpha /\
    U = 2 ^ alpha /\ V = 3 ^ alpha /\
    (forall k : Nat,
      let e := Int.floor (k * alpha)
      normalizedTwo k = U ^ k / 2 ^ e /\
      normalizedThree k = V ^ k / 3 ^ e) := by
  ...

/-- The automatic all-prime divisor of the balanced mixed determinant
cannot fill more than four fifths of its leading height. -/
theorem mixed_tropical_divisor_limsup_le_four_fifths
  (hprimitive : min (vTwo M) (vThree A) = 0) :
  limsup (fun T => log (tropicalDivisor T) / leadingHeight T)
  <= (4 : Real) / 5 := by
  ...
```

```

/-- The complete conditional grid has countermonotone/maximal transport
    ratio pi/4. -/
theorem full_grid_transport_ratio :
  countermonotoneEnergy / maximalEnergy = Real.pi / 4 := by
  ...

```

The actual APIs will need natural-valued denominator exponents rather than casts of negative integers, and the asymptotic statements will likely use `Tendsto` and `IsLittleO` rather than a custom `limsup`. The sketches are included to fix the mathematical interfaces, not Lean syntax.

44.4 The formalized finite core of the first tied layer

The reusable initial form of a determinant at a prime is now implemented. For a matrix over $\mathbb{Z}$ and an assignment minimum τ_p , reduce

$$p^{-\tau_p} \det(E_{ij}) \pmod{p}.$$

Only minimizing permutations survive. The modules `TropicalInitialForm`, `BlockFiberFactorization`, and `BeattyTropicalFiber` now prove, as purely finite identities, that the relevant minimizers preserve the equal- u or equal- v fibers, that the surviving sum is a product of smaller determinants, and that the next assignment layer begins at least $\lfloor \beta \rfloor$ later. The module `RowBlockLaplaceExpansion` extends this from the first fiber to every finite depth: after extracting the minimum power, its cutoff- K identity discards all row partitions of excess at least K modulo p^K and expresses every surviving coefficient as a signed product of square unit minors. This is an exact support decomposition, not a bound on the surviving signed sum. The companion `AssignmentExcessGeometry` proves that Beatty excess dominates $\lfloor \beta \rfloor$ times prefix-crossing energy and, for $N > 0$, bounds the number of partitions below cutoff K by $\binom{N+q-1}{q}$, with $q = \lfloor (K-1)/\lfloor \beta \rfloor \rfloor$. The finite valuation estimate is separated into `BlockVandermondeValuationBudget`. Combining its kernel-verified cube inequality with the elementary paper-level count and size bounds for the Beatty fibers gives

$$v_2(I_{2,T}) + v_3(I_{3,T}) = O(N^2) = o(T^{5/2}). \quad (219)$$

The optional classical two-logarithm input sharpens $O(N^2)$ to $O(N^{3/2}(\log N)^2)$, but the sharper exponent is not needed for the qualitative conclusion. Finally, `RowBlockUnitDivisibility` applies the unit-ordering divisor to every row-partition minor product and proves the additive scaled-block conclusion $p^{\tau+U} \mid \det$. Thus the reusable finite algebra and valuation inequality are formalized, while the displayed asymptotic specialization remains a short paper argument. The remaining mathematical target is no longer the first tied fiber; it is the valuation and cancellation of the low-excess truncated sums as their cutoff grows.

44.5 Certificate architecture for the new finite bound

The existing verified bound through 4096 should remain the trusted baseline until a new certificate is accepted. A practical route to 2^{27} is hierarchical:

1. partition $[1, 2^{27})$ into fixed-width chunks;
2. for each non-power-of-two m , store an integer nearest-value locator n_m and rational lower/upper enclosures proving either $m^\vartheta < n_m$ or $m^\vartheta > n_m$;

3. hash each chunk and prove a small checker correct in Lean;
4. combine the chunk theorems into one interval theorem.

The MPFR run in Section 40 is reconnaissance and a source of extremal cases; it is not the certificate itself.

45 Numerical sanity checks and reproducibility

None of the limiting constants in the determinant theorems depends on numerical evidence. The computations below were used to catch normalization errors, check convergence rates, and stress-test the closed forms before writing the proofs.

45.1 Balanced-block assignment constants

For each N , take the first N pairs (u, v) ordered by $u + \vartheta v$, pair them against N equally spaced row parameters, and solve the one-dimensional assignment problems by sorting. In the cross-coprime model, the structural places should tend to $3/10$ and the minimal all-prime three-vector model to $7/10$.

N	structural 2, 3 fraction	all-prime three-vector fraction
25	0.3224571184	0.7258906037
50	0.3135103379	0.7141968856
100	0.3050363664	0.7053583986
250	0.3024422068	0.7025311250
500	0.3012614628	0.7013136826
1000	0.3006780555	0.7006854844
2500	0.3002789586	0.7002894348
5000	0.3001354539	0.7001384826
Limit	$3/10$	$7/10$

Table 10: Finite assignment sums converging to the exact mixed-block constants.

The convergence from above is consistent with the $O(L^{-1}) = O(N^{-1/2})$ boundary error in triangular lattice counting. The table also exposes the error in the proposed $2/3$ ceiling: the finite all-prime ratio is already near 0.70 , and the exact limit is $7/10$.

45.2 Complete-grid transport ratio

For a representative hypothetical height parameter $\beta = 27.314159$, form the row lattice $n + \beta k \leq T$ and the matching low-weight column lattice. Normalize the countermonotone assignment by the largest-term assignment. The finite ratios are:

The slow boundary convergence is expected: both empirical triangles have $O(T)$ boundary points against $O(T^2)$ bulk points, and the matching weight scale also grows with T .

T	number of rows	transport ratio
50	73	0.7415113
100	237	0.7559915
200	841	0.7671191
400	3157	0.7750371
800	12210	0.7796781
Limit	—	$\pi/4 = 0.7853981634\dots$

Table 11: Finite full-grid countermonotone/maximal transport ratios.

45.3 Directed-rounding scan manifest

The scan in Section 40 used the literal C listing extracted from the attached LaTeX source. Its SHA-256 digest is

f888c0d89be9366f01e2d26f339dcbe446556d39d4c5dbcf8f579ec0d0052e1d.

This differs from the digest printed in the old report. To ensure the discrepancy was only a manifest error, the extracted source was first rerun on the old overlap interval $[29,360,128,33,554,432)$; it reproduced the old extremal records and reported zero failures.

The extension interval $[2^{26}, 2^{27})$ was split into sixteen chunks of width 2^{22} and run at 192 bits. The concatenated chunk log has SHA-256

b11843448cd8946442aba13758b5b666e4ac6eae50d2243fdbc32d83274cea8d.

The two new extremal records were replayed at 320 and 512 bits. A representative build and invocation sequence was:

```

1 cc -O3 -std=c17 mpfr_scan.c \
2   /lib/x86_64-linux-gnu/libmpfr.so.6 -lgmp -lm -o mpfr_scan
3
4 ./mpfr_scan 67108864 71303168 192
5 # repeat consecutive width-222 chunks through 134217728
6 sha256sum extension-through-2pow27.log
```

The scanner excludes powers of two from the tested population because they are the known solutions. Across the extension it checked 67,108,863 non-powers of two and reported zero failures.

45.4 Extremal replay records

The narrowest new lower-directed separation was

$$m = 132,833,765, \quad n = 7,501,348,219,569, \quad \text{directed margin } 5.9082259992879571 \times 10^{-21}.$$

The corresponding high-precision ordinary difference is approximately

$$m^\vartheta - n = 6.39397545334704 \times 10^{-8}.$$

The narrowest new upper-directed separation was

$$m = 129,799,345, \quad n = 7,231,571,292,851, \quad \text{directed margin } 8.3073231820333117 \times 10^{-22},$$

with

$$n + 1 - m^\vartheta = 8.666990435581896 \times 10^{-9}.$$

These decimal differences are explanatory only. The certification decision uses the MPFR outward-rounded interval endpoints and the scanner’s locator checks.


46 Status matrix and dependency audit

Statement	Status	Dependency or caveat
Exact primitive monoid $\mathcal{S} = \mathbb{N}_0 + \mathbb{N}_0\beta$ and candidate monoid	[kernel-verified Lean]	Inherited from the source report.
Primitive-depth condition $\min(v_2(M), v_3(A)) = 0$	[kernel-verified Lean]	Inherited; central to both universal ceilings.
Synchronized dyadic–triadic Sturmian recurrence	[paper proof in this report]	Elementary from the exact monoid and floor normalization.
Lowest-term denominator formulas and full-depth alternative	[paper proof in this report]	Unique factorization plus primitive depth.
Effective polynomial irrationality measure for fixed M	[classical/open background] plus [paper proof in this report]	Baker–Wustholz/Matveev; constants depend effectively on M .
Polynomial recurrence of continued-fraction denominators	[paper proof in this report]	Standard convergent inequality applied to the previous row.
Exact ordinary block valuations at all input-support primes	[paper proof in this report]	Unequal valuation of the two monomials in every difference.
Ordinary product fraction (131)	[paper proof in this report]	Exact floor sums; output endpoint errors are $O(T^2)$.
Universal ordinary ceiling $0.871049\dots$	[paper proof in this report]	Primitive depth only; not a bound on cancellation from new primes in differences.
Explicit triangular Gini norm	[paper proof in this report]	Direct integration.
Mixed all-prime tropical limit (154)	[paper proof in this report]	Rearrangement plus triangular Riemann sums.
Universal 4/5 mixed ceiling	[paper proof in this report]	Conservation, Gini triangle inequality, primitive depth.

Statement	Status	Dependency or caveat
Cross-coprime 7/10 and structural 3/10	[paper proof in this report]	Specialization $a = b = c = d = 0$.
Full-grid $\pi/4$ ceiling	[paper proof in this report]	Countermonotone and comonotone transport for the common quantile $\sqrt{t}$.
HCIZ determinant identity	[classical/open background]	Exact finite identity; the needed large- N inequality remains open.
First tied-layer fiber factorization	[paper proof in this report]	Exact initial-form determinant algebra in the cross-coprime branch.
First tied layer is $o(H_T)$	[classical/open background] plus [paper proof in this report]	Standard p -adic two-logarithm estimates; no claim about higher valuation layers.
Mixed cancellation-excess criterion	[paper proof in this report]	Immediate sufficient condition from integrality and the Gini deficit.
Software exclusion through 2^{27}	[software computation]	Complete MPFR directed-rounding scan; not a Lean proof.
Kernel exclusion through 4096	[kernel-verified Lean]	Remains the strongest formal finite theorem audited in the source report.
The Alaoglu–Erdős conjecture itself	[proposed target]	Equivalent to the unresolved special four-exponentials determinant.

46.1 Logical dependency graph

The new deductions can be summarized as



Baker–Matveev separation and the finite scan are parallel inputs: they sharpen the possible counterexample but do not enter the determinant ceiling proofs.

47 Conclusions

The Alaoglu–Erdős conjecture remains unproved. The continuation nevertheless changes the shape of the search in several concrete ways.

First, the exact conditional monoid has a compact symbolic normal form. After removing the integer part of the primitive generator, one irrational Sturmian word controls both the binary and ternary denominator jumps of two rational recurrences. This exposes a possible route through finite-state rigidity, but it also prevents a false shortcut: rational multipliers do not by themselves make the word automatic.

Second, ordinary transcendence theory gives more Diophantine separation than the source report’s elementary argument records. A fixed hypothetical generator is a ratio of two logarithms of integers, so Baker–Wustholz/Matveev yields an effective polynomial irrationality measure. Continued-fraction denominators therefore grow at most polynomially from one stage to the next. This is useful quantitative structure, but it does not conflict with the $O(T)$ exact population of a dyadic block.

Third, the most direct p -adic determinant plan can now be audited with exact constants. For ordinary block Vandermondes, every valuation at every prime already supporting the inputs is accounted for, and primitive depth leaves at least 12.8950...% of the leading archimedean budget uncovered. For the balanced mixed determinant, all primes collapse into a finite family of slope vectors summing to zero. The triangular Gini norm then proves that entrywise divisibility misses at least 20% universally, at least 30% in the cross-coprime branch, and at least 70% if one uses only the structural places 2, 3 there. The complete two-dimensional grid improves the geometry but still leaves a termwise ceiling of $\pi/4$.

These are no-go theorems for particular mechanisms, not evidence that the conjecture is false. They identify the exact remaining mechanisms:

1. a leading-order cascade of cancellations among many equal-valuation Leibniz terms;

2. an archimedean determinant estimate substantially smaller than its largest-term bound, most naturally through the HCIZ free energy;
3. a genuinely mixed real/ p -adic construction that turns $\log M \log 3 = \log A \log 2$ into deep congruences; or
4. a rigidity theorem for the synchronized rational Sturmian system.

The first cancellation layer is provably too small, so any successful p -adic proof must propagate through many valuation layers rather than extracting one convenient Vandermonde factor.

Finally, the directed-rounding computation was extended by one full binary exponent. At the software-trust level, no non-power-of-two candidate occurs below 2^{27} , so a counterexample would satisfy $x > 27$. The formal Lean boundary remains $x \geq 12$ until a generated certificate is checked by the kernel. The old scanner-source hash printed in the source report is incorrect; the literal source, overlap replay, new log hash, and high-precision extremal replays are recorded explicitly in this document.

Bottom line. The conjecture is still open, but the continuation produces new exact mathematics rather than only a list of ideas: a synchronized Sturmian equivalence; effective polynomial separation; closed all-prime Vandermonde asymptotics; a triangular Gini norm governing mixed tropical divisibility; universal $4/5$, $7/10$, $3/10$, and $\pi/4$ ceilings; a first-layer cancellation theorem; exact sufficient thresholds for future work; and a reproducible software exclusion through 2^{27} . The next decisive calculation is the triangular HCIZ free energy or, on the arithmetic side, a proof of a genuinely leading tropical cancellation cascade.

48 Second continuation: scope and claim ledger

The problem is

This document is a sequel rather than a replacement for the preceding parts of this report. It restates every inherited fact used below, because several later arguments depend sensitively on the exact quantifiers and on whether a quantity is integral or merely rational. In particular, the slogan “one counterexample generates a grid” is proved here as an exact equality of monoids, not used as a heuristic.

No complete proof of Conjecture 2.1 is claimed. The strongest new unconditional outputs of this continuation are:

1. a streamlined proof of the exact free-monoid structure, including the denominator-elimination step that is easy to understate;
2. a universal N^2 -scale prime-divisibility lower bound for exact-grid interpolation determinants;
3. the triangular determinant inequality of Theorem 53.1;
4. a rational-basis, product-formula extension that quantifies exactly how denominators defeat naive lattice reduction;
5. additional shared-prime, congruence, counting, and formalization consequences.

48.1 Claim ledger

Claim	Status	Main input
A counterexample is positive and transcendental	proved	elementary reductions; Gelfond–Schneider
All solutions lie in a rational affine line through 1 and one nonintegral solution	proved	Six Exponentials Theorem
The complete solution set is $\mathbb{N}_0 + \mathbb{N}_0\beta$	proved	valuations and minimality
A counterexample has finite effective irrationality measure	proved	linear forms in two logarithms
Universal square-grid divisor $6^{(1+\beta)L_m}$	proved	tropical determinant valuation
Triangular inequality $g(\sqrt{\beta \log 2 \log 3}) \geq 0$	proved	HCIZ bound, assignment asymptotics, Baker uniformity
Rational-basis denominator inequality	proved	product formula and signed assignment bounds
A complete contradiction for all β	not obtained	requires further N^2 -scale arithmetic gain

The elementary reductions, the affine rank bound, the primitive least counterexample and the exact free-monoid theorem restated at this point in the source document are already established earlier in this report, and in kernel-verified form: see Section 5 and the modules `LeastSolution`, `PrimitiveGenerator` and `RationalPowerIndex`. They are not repeated here. What follows is the new material of the second continuation, which takes the exact grid as given and asks how much arithmetic a determinant built on it can extract.

49 Effective Diophantine and counting consequences

The two logarithmic descriptions of β imply more than transcendence. Because A is not a power of 2, the linear form $q \log A - p \log 2$ is nonzero whenever $(p, q) \neq (0, 0)$. Effective lower bounds for linear forms in logarithms therefore give constants $c > 0$ and $\tau > 0$, effectively computable from A , such that

$$\|q\beta\| \geq cq^{-\tau} \quad (q \geq 1). \quad (220)$$

The analogous representation $\beta = \log B / \log 3$ gives another such bound. No useful small absolute exponent is asserted here; standard explicit bounds are enormous. The qualitative conclusion is nevertheless important.

Corollary 49.1 (Finite irrationality measure). *A counterexample β is not a Liouville number. Its irrationality measure is finite and effectively bounded in terms of either A or B .*

The exact monoid also gives an exact counting formula. Let

$$N_{\mathcal{S}}(T) := \#(\mathcal{S} \cap [0, T]).$$

Then, with $K = \lfloor T/\beta \rfloor$,

$$N_{\mathcal{S}}(T) = \sum_{k=0}^K (\lfloor T - k\beta \rfloor + 1). \quad (221)$$

Proposition 49.2 (Quadratic counting with a power saving). *There is an effective $\delta > 0$ such that*

$$N_{\mathcal{S}}(T) = \frac{T^2}{2\beta} + \frac{\beta + 1}{2\beta}T + O_{\beta}(T^{1-\delta}).$$

Proof sketch. Equation (220), the Erdős–Turán inequality, and the geometric-series estimate for $\sum_{k \leq K} e^{2\pi i h k \beta}$ give a power saving for the discrepancy of the rotation sequence $\{k\beta\}$. In particular,

$$\sum_{k=0}^K \{T - k\beta\} = \frac{K+1}{2} + O(K^{1-\delta})$$

for some effective $\delta > 0$. Insert this in Equation (221). Writing $K = T/\beta - \{T/\beta\}$, the apparently oscillatory terms of order T cancel, leaving the displayed linear coefficient. $\square$

The same argument gives quantitative equidistribution of the dense orbit

$$(2^{\{k\beta\}}, 3^{\{k\beta\}}) = \left(\frac{A^k}{2^{\lfloor k\beta \rfloor}}, \frac{B^k}{3^{\lfloor k\beta \rfloor}} \right).$$

Thus a counterexample would generate a dense family of rational points on the arc $t \mapsto (2^t, 3^t)$, with pure base-power denominators. Their heights grow exponentially in k , so this observation alone is far below the density needed for Pila–Wilkie-type counting to force algebraicity; it is nevertheless a useful reformulation for S -unit methods.

50 Valuation-lattice saturation and an exact duality

The denominator argument has a useful invariant formulation. It also has a symmetric consequence that is easy to miss if one works only with the least counterexample β .

Write

$$a_p = v_p(A), \quad b_p = v_p(B),$$

and form the two-column integer matrix whose rows are

$$(\mathbf{1}_{p=2}, a_p) \quad \text{and} \quad (\mathbf{1}_{p=3}, b_p) \quad (p \text{ prime}).$$

Only finitely many entries in the second column are nonzero. Let $\mathcal{L}$ be the rank-two lattice generated by the columns in the direct sum of all row copies of $\mathbb{Z}$.

Theorem 50.1 (Saturated valuation lattice). *The lattice $\mathcal{L}$ is primitive (saturated):*

$$(\mathcal{L} \otimes_{\mathbb{Z}} \mathbb{Q}) \cap \bigoplus_p (\mathbb{Z} \oplus \mathbb{Z}) = \mathcal{L}.$$

Equivalently, the greatest common divisor of all maximal minors of the matrix is 1. In the present sparse matrix this is exactly

$$\gcd(\{a_p : p \neq 2\} \cup \{b_p : p \neq 3\} \cup \{a_2 - b_3\}) = 1.$$

Proof. The nonzero maximal minors are, up to sign, exactly the displayed a_p , b_p , and $a_2 - b_3$: the first column is nonzero only in the row belonging to $p = 2$ in the first valuation block and the row belonging to $p = 3$ in the second. It remains to prove that their gcd is 1.

If a number $d > 1$ divided all these minors, choose $t \in \{0, \dots, d-1\}$ with $t \equiv -a_2 \pmod{d}$. The congruence $a_2 \equiv b_3 \pmod{d}$ then gives $t \equiv -b_3 \pmod{d}$, while every other valuation of A or B is already divisible by d . Consequently both $2^t A$ and $3^t B$ are perfect d th powers in $\mathbb{N}$. Hence

$$\frac{\beta + t}{d} \in S.$$

This number is nonintegral, and it is strictly between 0 and β because $\beta > 1$ and $0 \leq t \leq d-1$, contradicting the choice of β . Thus the gcd of the maximal minors is 1. The rank-two Smith normal form criterion then says exactly that the column lattice is saturated. $\square$

A particularly concrete form of the theorem records exactly how extraction of common roots behaves throughout the solution monoid.

Corollary 50.2 (Perfect-power content is preserved). *For $u, v \in \mathbb{N}_0$, not both zero, define*

$$X_{u,v} = 2^u A^v, \quad Y_{u,v} = 3^u B^v.$$

Then

$$\gcd(\{v_p(X_{u,v}) : p \text{ prime}\} \cup \{v_p(Y_{u,v}) : p \text{ prime}\}) = \gcd(u, v).$$

Equivalently, for every $d \geq 1$, the two integers $X_{u,v}$ and $Y_{u,v}$ are simultaneously perfect d th powers if and only if $d \mid u$ and $d \mid v$.

Proof. Every divisor of $\gcd(u, v)$ plainly divides all displayed valuations. In the other direction, if d divides all of them, then $X_{u,v} = C^d$ and $Y_{u,v} = D^d$ for positive integers C, D . Thus $(u + v\beta)/d \in S$. Uniqueness of the representation in the irrational basis $(1, \beta)$ gives $u/d, v/d \in \mathbb{N}_0$. $\square$

For example, A and B cannot themselves be simultaneous proper powers; and for every $t \geq 0$, neither pair

$$(2^t A, 3^t B), \quad (2A^t, 3B^t)$$

can consist of simultaneous proper powers of the same degree. These are not independent ad hoc restrictions: all of them are specializations of one saturated rank-two valuation lattice.

There is also an exact reciprocal description. Put

$$S_{A,B} = \{y \in \mathbb{R} : A^y \in \mathbb{Z}, B^y \in \mathbb{Z}\}.$$

Theorem 50.3 (Reciprocal solution monoid). *Under the counterexample hypothesis,*

$$S_{A,B} = \mathbb{N}_0 + \frac{1}{\beta}\mathbb{N}_0.$$

In particular $1/\beta$ is the least positive nonintegral element of $S_{A,B}$.

Proof. If $y \in S_{A,B}$, then $x = \beta y$ satisfies

$$2^x = A^y \in \mathbb{Z}, \quad 3^x = B^y \in \mathbb{Z}.$$

Hence $x = n + k\beta$ with $n, k \in \mathbb{N}_0$, and $y = k + n/\beta$. Conversely,

$$A^{k+n/\beta} = A^k 2^n, \quad B^{k+n/\beta} = B^k 3^n$$

are integers for all $n, k \geq 0$. $\square$

This reciprocal monoid shows that the two columns of the valuation matrix may be interchanged without losing arithmetic information: the original and dual root-extraction problems are the same saturated lattice viewed in opposite bases.

51 Exact-grid interpolation determinants

The free monoid of the exact primitive-generator theorem (Section 5) supplies two exact rank-two grids: one in the logarithms of the bases and one in the solution exponents. Retaining both coordinates is the central step of this continuation.

For finite sets $\mathcal{R}, \mathcal{C} \subset \mathbb{N}_0^2$ with $|\mathcal{R}| = |\mathcal{C}| = N$, define

$$\begin{aligned} u_{a,b} &:= a\alpha + b\gamma, & (a,b) &\in \mathcal{R}, \\ v_{n,k} &:= n + k\beta, & (n,k) &\in \mathcal{C}, \end{aligned}$$

and the matrix

$$M_{(a,b),(n,k)} := \exp(u_{a,b}v_{n,k}) = 2^{an}3^{bn}A^{ak}B^{bk}. \quad (222)$$

Every entry is a positive integer.

51.1 Nonvanishing and total positivity

Lemma 51.1 (Strict positivity). *Order the u -values and the v -values increasingly. Then*

$$D(\mathcal{R}, \mathcal{C}) := \det(\exp(u_i v_j))_{i,j=1}^N > 0.$$

In particular it is a positive integer and hence at least 1.

Proof. The u -values are distinct because 2 and 3 are multiplicatively independent, and the v -values are distinct because β is irrational. The kernel $K(u, v) = e^{uv}$ is strictly totally positive on $\mathbb{R} \times \mathbb{R}$. Equivalently, the functions $e^{v_1 u}, \dots, e^{v_N u}$ with $v_1 < \dots < v_N$ form an extended complete Chebyshev system: their Wronskian is

$$e^{u(v_1 + \dots + v_N)} \prod_{i < j} (v_j - v_i) > 0.$$

The generalized mean-value theorem for Chebyshev systems gives the desired sign. $\square$

51.2 Archimedean upper bound

Write $\Delta(u) = \prod_{i < j} (u_j - u_i)$ and similarly for v , and let $G(N+1) = \prod_{j=1}^{N-1} j!$ be the relevant Barnes product.

Lemma 51.2 (HCIZ determinant inequality). *For increasingly ordered real vectors u, v ,*

$$0 < \det(e^{u_i v_j}) \leq \frac{\Delta(u) \Delta(v)}{G(N+1)} \exp\left(\sum_{i=1}^N u_i v_i\right). \quad (223)$$

Proof. The Harish–Chandra–Itzykson–Zuber formula with normalized Haar measure is

$$\int_{U(N)} e^{\text{tr}(\text{diag}(u)U \text{diag}(v)U^*)} dU = G(N+1) \frac{\det(e^{u_i v_j})}{\Delta(u) \Delta(v)}.$$

Von Neumann’s trace inequality bounds the exponent in the integral above by $\sum_i u_i v_i$. $\square$

The factor $G(N+1)$ is crucial. When both node sets come from two-dimensional lattice regions, the leading $N^2 \log N$ contributions of the two Vandermonde products cancel the $N^2 \log N$ contribution of $G(N+1)$ exactly. The problem therefore lives at the next, N^2 -constant, scale. This is the determinant manifestation of the critical nature of four exponentials.

51.3 Finite-prime lower bound as an assignment problem

For a prime p , the valuation of an entry in (222) is

$$w_p((a, b), (n, k)) = \mathbf{1}_{p=2} a n + \mathbf{1}_{p=3} b n + a_p a k + b_p b k. \quad (224)$$

Lemma 51.3 (Tropical determinant bound). *For every prime p ,*

$$v_p(D(\mathcal{R}, \mathcal{C})) \geq \min_{\sigma \in S_N} \sum_{i=1}^N w_p(i, \sigma(i)).$$

Proof. Each Leibniz term has the displayed summed valuation. The nonarchimedean triangle inequality says that the valuation of their signed sum is at least the minimum of their valuations. $\square$

This converts divisibility into a finite optimal-transport problem. Even a coarse componentwise use of it already produces an N^2 -scale divisor.

52 Square grids: a universal divisor and the critical cancellation

Take $\mathcal{R} = \mathcal{C} = I_m := \{0, \dots, m-1\}^2$ and $N = m^2$. Denote the corresponding determinant by D_m .

Let L_m be the minimum of $\sum_i a_i n_{\sigma(i)}$ over bijections between the two copies of I_m . The multiset of first coordinates consists of each integer $0, \dots, m-1$ repeated m times. The rearrangement inequality gives

$$L_m = m \sum_{j=0}^{m-1} j(m-1-j) = \frac{m^2(m-1)(m-2)}{6}. \quad (225)$$

The same value applies to each of the four pairings an, bn, ak, bk .

Theorem 52.1 (Universal square-grid divisor). *For every $m \geq 1$,*

$$\log D_m \geq L_m(\log 2 + \log 3 + \log A + \log B) = L_m(1 + \beta) \log 6. \quad (226)$$

Equivalently, the integer D_m is at least $6^{(1+\beta)L_m}$ in logarithmic size.

Proof. For a fixed prime, the minimum of a sum of nonnegative cost components is at least the sum of their separate minima. Apply Lemma 51.3 to (224); each present component has minimum L_m . Sum the resulting lower bounds with weights $\log p$ and use $\sum_p a_p \log p = \log A$ and $\sum_p b_p \log p = \log B$. $\square$

The estimate is much stronger than the bare fact $D_m \geq 1$. Nevertheless, a balanced rectangular asymptotic does not yet contradict it. Let $z = \sqrt{\alpha\gamma\beta}$. After choosing rectangle aspect ratios so that the two contributions to each projected logarithm are equal, the normalized node distribution is that of $Z = (X + Y)/2$ for independent $X, Y \sim \text{Unif}[0, 1]$. One has

$$\mathbb{E}Z^2 = \frac{7}{24}, \quad \iint \log |s - t| d\text{Law}(Z)(s) d\text{Law}(Z)(t) = \frac{1}{3} \log 2 - \frac{25}{12}.$$

Combining these with Lemma 51.2 and Theorem 52.1 yields the necessary rectangular inequality

$$h(z) := \frac{1}{2} \log(4z) + \frac{1}{3} \log 2 - \frac{4}{3} + \frac{z}{2} \geq 0. \quad (227)$$

For the range allowed even by $A \geq 3$, this inequality is already positive. The square-grid divisor is therefore real progress in the arithmetic estimate, but rectangular geometry spends too much archimedean size.

53 Triangular grids and a new exclusion inequality

The natural way to reduce the archimedean cost is to take the first N nodes rather than a rectangular block.

Let $\mathcal{R}_N$ consist of the N pairs $(a, b) \in \mathbb{N}_0^2$ with the smallest values of $u = a\alpha + b\gamma$, and let $\mathcal{C}_N$ consist of the N pairs $(n, k) \in \mathbb{N}_0^2$ with the smallest values of $v = n + k\beta$. Ties do not occur. Let D_N be the positive integer determinant from these sets.

Define

$$U := \sqrt{2\alpha\gamma}, \quad V := \sqrt{2\beta}, \quad z := \sqrt{\alpha\gamma\beta},$$

so $UV = 2z$.

53.1 Limiting distributions

Elementary lattice-point counting gives

$$\#\{(a, b) \in \mathbb{N}_0^2 : a\alpha + b\gamma \leq T\} = \frac{T^2}{2\alpha\gamma} + O(T),$$

and similarly for $n + k\beta$. Thus the N th thresholds are asymptotic to $U\sqrt{N}$ and $V\sqrt{N}$. After division by $\sqrt{N}$, the projected nodes converge to the distributions

$$d\mu_U(t) = \frac{2t}{U^2} dt \quad (0 \leq t \leq U), \quad d\nu_V(t) = \frac{2t}{V^2} dt \quad (0 \leq t \leq V).$$

Their quantile functions are

$$Q_{\mu_U}(s) = U\sqrt{s}, \quad Q_{\nu_V}(s) = V\sqrt{s}.$$

The logarithmic energies are

$$\mathcal{E}(\mu_U) = \log U - \frac{7}{4}, \quad \mathcal{E}(\nu_V) = \log V - \frac{7}{4}. \quad (228)$$

The finite irrationality type supplied by linear forms in logarithms ensures that exceptionally close projected lattice points do not spoil convergence of the logarithmic energies; details are given in Appendix M.

53.2 The archimedean constant

The Barnes asymptotic is

$$\log G(N+1) = \frac{N^2}{2} \log N - \frac{3N^2}{4} + o(N^2).$$

Using (228), the two Vandermonde terms in Lemma 51.2, and comonotone coupling for the trace term, gives

$$\limsup_{N \rightarrow \infty} \frac{\log D_N}{N^2} \leq \frac{1}{2} \log(UV) - 1 + \frac{UV}{2}. \quad (229)$$

Indeed,

$$\int_0^1 Q_{\mu_U}(s) Q_{\nu_V}(s) ds = UV \int_0^1 s ds = \frac{UV}{2}.$$

53.3 The finite-prime constant

Uniform measure on the row triangle can be written as

$$a = \frac{U}{\alpha}X, \quad b = \frac{U}{\gamma}Y, \quad X, Y \geq 0, \quad X + Y \leq 1,$$

where (X, Y) is uniform on the standard simplex. Each coordinate has density $2(1 - x)$ and quantile

$$q(s) = 1 - \sqrt{1 - s}.$$

The column coordinates have the corresponding scales $n = VX'$ and $k = (V/\beta)Y'$. The counter-monotone assignment constant is

$$J := \int_0^1 q(s)q(1 - s) ds = \frac{\pi}{8} - \frac{1}{3}. \quad (230)$$

Consequently each of the four logarithmically weighted components $an \log 2$, $bn \log 3$, $ak \log A$, and $bk \log B$ contributes asymptotically $UVJN^2$ to the componentwise valuation lower bound. Therefore

$$\liminf_{N \rightarrow \infty} \frac{\log D_N}{N^2} \geq 4UVJ. \quad (231)$$

Theorem 53.1 (Triangular determinant inequality). *Every nonintegral counterexample β satisfies*

$$g(\sqrt{\beta \log 2 \log 3}) \geq 0, \quad g(z) = \frac{1}{2} \log(2z) - 1 + \left(\frac{11}{3} - \pi\right)z. \quad (232)$$

Proof. Subtract the lower bound (231) from the upper bound (229). Since $UV = 2z$ and $4J = \pi/2 - 4/3$, the resulting necessary inequality is exactly (232). $\square$

The derivative $g'(z) = 1/(2z) + 11/3 - \pi$ is positive for $z > 0$, so g has a unique positive zero

$$z_0 = 1.12895010749713, \quad \beta_\Delta := \frac{z_0^2}{\log 2 \log 3} = 1.67370758736671.$$

At the first arithmetically possible nonintegral value $\beta = \log_2 3$, one has $z = \log 3$ and

$$g(\log 3) = -0.029549732685145 < 0.$$

The numerical consequence of Theorem 53.1 is weaker than what is already known: it gives $\beta \geq \beta_\Delta = 1.6737\dots$, whereas the kernel-verified finite check of Section 25 gives $\beta \geq 12$ and the directed-rounding scan reaches 2^{27} . The inequality is recorded for its method, not for its bound: it is the first exclusion obtained from a uniform asymptotic determinant estimate rather than from a per-value certificate, and it exposes the arithmetic constant that would have to improve in order to reach all A at once.

53.4 Numerical anatomy of the two grid geometries

The table compares the triangular gap g with the balanced-rectangle gap h from (227). Negative values would contradict a counterexample at that β .

β	$z = \sqrt{\alpha\gamma\beta}$	$g(z)$	$h(z)$
1	0.8726396796	-0.2633422644	-0.04093352603
$\log_2 3$	1.098612289	-0.02954973269	0.1871929656
β_Δ	1.128950107	-1.054219794e-81	0.2159820074
2	1.23409887	0.09973735836	0.3130828643
$\log_2 5$	1.329717364	0.1872568287	0.3982047952
3	1.511456262	0.3467367942	0.5531278372
5	1.951281643	0.7053840783	0.900746934
10	2.759528964	1.303060538	1.47815739

The triangular region improves the constant substantially, but g is increasing and becomes comfortably positive in the range $\beta \geq \log_2 5$. To continue, one must either improve the finite-prime lower bound, find a better variational region, or enrich the determinant.

54 Rational changes of basis and the product formula

A natural reaction to the positive gap is to replace 3 by a ratio such as $3/2$, whose logarithm is smaller. Ignoring denominators makes this look like an immediate proof. The product formula shows exactly why that argument is false and, more usefully, turns it into quantitative denominator constraints.

54.1 A signed arithmetic functional

For a rational $q > 1$, write it in lowest terms as $q = \text{num}(q)/\text{den}(q)$ and define

$$\rho(q) := \frac{J \log \text{num}(q) - K \log \text{den}(q)}{\log q}, \quad J = \frac{\pi}{8} - \frac{1}{3}, \quad K = \frac{1}{6}. \quad (233)$$

The constant J is the countermonotone coordinate cost in a simplex, while K is the comonotone cost, since a simplex coordinate has second moment $1/6$. If q is an integer, then $\rho(q) = J$. A denominator contributes with K , not J , because a negative valuation reverses the relevant assignment bound.

Theorem 54.1 (Adelic triangular inequality). *Let $r_1, r_2 > 1$ be multiplicatively independent positive rationals and suppose $R_i = r_i^\beta \in \mathbb{Q}_{>1}$ for $i = 1, 2$. Put*

$$z_r := \sqrt{\beta \log r_1 \log r_2}.$$

Then the existence of these four rational exponentials implies

$$\frac{1}{2} \log(2z_r) - 1 + z_r - 2z_r(\rho(r_1) + \rho(r_2) + \rho(R_1) + \rho(R_2)) \geq 0. \quad (234)$$

Equivalently, if

$$\Delta_{\text{den}} := \sum_{q \in \{r_1, r_2, R_1, R_2\}} \frac{\log \text{den}(q)}{\log q},$$

then

$$g(z_r) + 2z_r \left(\frac{1}{2} - \frac{\pi}{8} \right) \Delta_{\text{den}} \geq 0. \quad (235)$$

Proof. Use the same first- N triangular sets for the positive logarithms $\log r_1, \log r_2$ and for $1, \beta$. The determinant entries are now rational. For a fixed prime, split each signed valuation coefficient into its positive and negative parts. In the tropical determinant bound, a positive coefficient is multiplied by the countermonotone minimum J , while a negative coefficient is multiplied by the comonotone maximum K . Summing over primes and using the product formula for the nonzero rational determinant gives the lower constant

$$2z_r(\rho(r_1) + \rho(r_2) + \rho(R_1) + \rho(R_2)).$$

The HCIZ upper constant remains $\frac{1}{2} \log(2z_r) - 1 + z_r$. This proves (234). Finally,

$$\rho(q) = J - (K - J) \frac{\log \text{den}(q)}{\log q},$$

and substitution gives (235). $\square$

54.2 All unimodular bases of $\langle 2, 3 \rangle$

Let

$$M = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \in \text{GL}_2(\mathbb{Z})$$

and suppose

$$\lambda_1 = a\alpha + b\gamma > 0, \quad \lambda_2 = c\alpha + d\gamma > 0.$$

Set

$$r_1 = 2^a 3^b, \quad r_2 = 2^c 3^d, \quad R_1 = A^a B^b, \quad R_2 = A^c B^d.$$

Then Theorem 54.1 applies. This supplies an infinite family of necessary inequalities indexed by unimodular changes of basis. Continued-fraction vectors make one λ_i small, but the corresponding rational base has a large denominator; Equation (235) measures the precise price.

Proposition 54.2 (Optimality of the integral basis). *Among unimodular bases for which both transformed bases are integers, the product $\lambda_1 \lambda_2$ is at least $\alpha\gamma$. Equality occurs only for the original basis $(2, 3)$ up to permutation. Likewise, among unimodular bases of the solution lattice whose two transformed exponents remain in $\mathcal{S}$, the product of the positive exponent generators is at least β .*

Proof. If $2^a 3^b$ is an integer, then $a, b \geq 0$. Thus an integral unimodular basis has nonnegative entries. Expanding,

$$(a\alpha + b\gamma)(c\alpha + d\gamma) \geq (ad + bc)\alpha\gamma \geq \alpha\gamma,$$

because $|ad - bc| = 1$. Equality forces the standard basis up to permutation. For the exponent lattice, the exact primitive-generator theorem says that a transformed exponent is simultaneously integral at bases 2 and 3 exactly when its coefficients in $(1, \beta)$ are nonnegative. The same argument applies. $\square$

This proposition explains why a denominator-free lattice reduction cannot improve the triangular inequality: the exact monoid has already selected the smallest integral basis.

54.3 The basis $(2, 3/2)$ and a gcd constraint

Take

$$r_1 = 2, \quad r_2 = \frac{3}{2}, \quad R_1 = A, \quad R_2 = \frac{B}{A}.$$

Let $d = A/\gcd(A, B)$ be the denominator of B/A in lowest terms and put $\delta = \log(3/2)$. Then

$$\Delta_{\text{den}} = \frac{\alpha}{\delta} + \frac{\log d}{\beta\delta}. \quad (236)$$

Theorem 54.1 therefore gives a rigorous lower bound on d , whenever the right side obtained by solving (235) is positive. In words: if $3/2$ substantially improves the archimedean configuration, then B/A must carry enough denominator to pay for that improvement. A counterexample cannot have excessive divisibility $A \mid B$ in the range where the reduced-basis gap is negative.

For orientation, the next table substitutes the hypothetical value $\beta = \log_2 A$. The column $d_{\min}$ is the least integer denominator not excluded by the inequality alone; it is not a claim that the corresponding B exists.

A	$\log_2 A$	z_r	gap before $\log d$	$d_{\min}$
3	1.5849625	0.667419621	-0.260297461	4
5	2.32192809	0.807818616	-0.0396109308	2
6	2.5849625	0.852347316	0.0269341919	1
7	2.80735492	0.88825597	0.0795954278	1
9	3.169925	0.94387388	0.159569469	1
10	3.32192809	0.966239056	0.191227186	1
11	3.45943162	0.986033907	0.219022705	1
12	3.5849625	1.00376439	0.243748126	1
13	3.70043972	1.01980266	0.265979188	1
14	3.80735492	1.03443012	0.286146736	1
15	3.9068906	1.04786443	0.304581094	1
17	4.08746284	1.07180649	0.337231592	1
18	4.169925	1.08256404	0.351820061	1
19	4.24792751	1.09264233	0.36544257	1
20	4.32192809	1.10211837	0.378212214	1
21	4.39231742	1.11105698	0.39022374	1
22	4.45943162	1.11951323	0.401557272	1
23	4.52356196	1.12753426	0.412281139	1

A second useful basis is

$$r_1 = \frac{3}{2}, \quad r_2 = \frac{4}{3}, \quad R_1 = \frac{B}{A}, \quad R_2 = \frac{A^2}{B}.$$

Its two logarithms are smaller, but Equation (235) now charges the denominators of all four displayed ratios. It therefore constrains the possibility that B/A and A^2/B are simultaneously close to integers. More generally, continued-fraction bases produce a network of constraints on

$$\frac{B^q}{A^p} \quad \text{when} \quad q \log 3 - p \log 2 \text{ is small.}$$

This is a concrete future route: prove that the denominator requirements for a carefully chosen finite family of bases are incompatible with the prime valuation vectors of A and B .

55 Joint prime transport and shared-factor gains

The componentwise lower bounds used above are intentionally conservative. At a fixed prime, a single permutation must handle all valuation components at once. If a prime occurs in more than one generator, the separate countermonotone assignments can conflict, making the true minimum strictly larger.

For the original triangular determinant, define

$$\mathcal{T}_p(N) := \frac{1}{N^2} \min_{\sigma \in S_N} \sum_i w_p(i, \sigma(i)).$$

Then the strongest direct product-formula lower bound is

$$\liminf_{N \rightarrow \infty} \frac{\log D_N}{N^2} \geq \sum_p (\log p) \liminf_{N \rightarrow \infty} \mathcal{T}_p(N).$$

The baseline $4UVJ$ replaces each $\mathcal{T}_p$ by the sum of four independent one-coordinate minima.

A simple rigorous gain estimate is available for a prime $p \notin \{2, 3\}$ that divides both A and B . Put $r = v_p(A)$, $s = v_p(B)$ and

$$R = \frac{r}{\alpha}, \quad S = \frac{s}{\gamma}.$$

On the normalized row simplex, the relevant scalar is $RX + SY$, while the column coordinate k has the marginal law of a simplex coordinate. The countermonotone cost is at least the product of the means minus the product of the standard deviations. Since

$$\mathbb{E}X = \mathbb{E}Y = \frac{1}{3}, \quad \text{Var } X = \text{Var } Y = \frac{1}{18}, \quad \text{Cov}(X, Y) = -\frac{1}{36},$$

one obtains

$$\liminf \mathcal{T}_p(N) \geq \frac{UV}{\beta} \left(\frac{R+S}{9} - \frac{\sqrt{R^2 + S^2 - RS}}{18} \right). \quad (237)$$

The componentwise baseline for the same two terms is

$$\frac{UV}{\beta} J(R+S).$$

Taking the larger of these two lower bounds gives a fully explicit shared-prime improvement. Summed over common prime divisors of A and B , it yields a sufficient contradiction criterion whenever the gain exceeds the positive gap $g(z)$.

The estimate is strongest when r/α and s/γ are comparable and vanishes continuously as one valuation goes to zero. It therefore does not supply a universal proof: the worst arithmetic configuration may have nearly disjoint prime support. But it turns “a large gcd should help” into a quantitative theorem and identifies exact valuation patterns worth targeting.

56 Universal congruence divisibility is lower order

There is also divisibility at primes not dividing $6AB$. Modulo such a prime p , a row of the matrix depends on (a, b) only modulo $p - 1$, by Fermat’s little theorem. Hence there are at most $(p - 1)^2$ distinct rows modulo p . For an $N \times N$ determinant,

$$v_p(D_N) \geq N - \text{rank}_{\mathbb{F}_p}(M) \geq N - (p - 1)^2$$

whenever the right side is positive. Summing over primes $p \leq \sqrt{N}$ and using the prime number theorem gives the universal extra estimate

$$\sum_{\substack{p \leq \sqrt{N} \\ p \nmid 6AB}} (N - (p - 1)^2) \log p = \left(\frac{2}{3} + o(1) \right) N^{3/2}. \quad (238)$$

Higher prime powers and column congruences can improve the lower-order term. They do not alter the N^2 constant in Theorem 53.1. This provides a useful barrier result: congruence collisions that depend only on finite-group periodicity are too small by a factor of $\sqrt{N}$. A complete determinant proof needs counterexample-specific valuation structure or a different determinant geometry.

57 Why several tempting routes still stop short

57.1 Continued fractions and integer gaps

For integers $p, q > 0$,

$$\frac{A^q}{2^p} = 2^{q\beta - p}, \quad \frac{B^q}{3^p} = 3^{q\beta - p}.$$

When p/q is a convergent to β , both ratios are close to 1. Since the cross-multiplied differences are nonzero integers, one obtains exponential lower bounds such as

$$|q\beta - p| \gg 2^{-p}.$$

These are far weaker than the polynomial lower bounds already supplied by linear forms in logarithms. Conversely, Baker’s polynomial lower bounds do not contradict the q^{-2} approximations supplied by continued fractions. The two representations of β give proportional linear forms, not two independent small quantities.

57.2 Catalan, abc , and small differences

The equations $A^q - 2^p = \pm 1$ and $B^q - 3^p = \pm 1$ would be highly constrained, but continued fractions only make the relative difference small, not the absolute difference. The absolute scale is exponential in q . Standard abc heuristics do not bridge that exponential gap without an additional source of unusually good approximation.

57.3 Counting rational points

The orbit of fractional parts gives $O(\log H)$ rational points of height at most H on a fixed transcendental arc. This is compatible with every known subpolynomial upper bound for rational points on definable transcendental curves. A counting theorem would need to exploit the pure 2-power and 3-power denominator restrictions, not merely rationality.

57.4 The four-exponentials boundary

If x is irrational, the pairs $(\log 2, \log 3)$ and $(1, x)$ are both $\mathbb{Q}$ -linearly independent, while all four exponentials

$$2, \quad 3, \quad 2^x, \quad 3^x$$

are algebraic. Thus the Four Exponentials Conjecture would prove Conjecture 2.1 immediately. The exact cancellation of the $N^2 \log N$ terms in our determinant estimates is the quantitative shadow of being exactly at this boundary. Six Exponentials is sufficient to collapse all possible exceptions to one affine direction, but not to remove that last direction.

58 A variational reformulation of the remaining determinant problem

The preceding calculations extend to scaled lattice regions $R, C \subset \mathbb{R}_{\geq 0}^2$ of equal area. Push forward uniform measure on R by $(a, b) \mapsto a\alpha + b\gamma$, and uniform measure on C by $(n, k) \mapsto n + k\beta$. The HCIZ side is determined by their logarithmic energies and comonotone transport. The finite-prime side is a sum of countermonotone or, more accurately, joint optimal-transport costs for the valuation bilinear forms (224).

This suggests the following concrete variational program.

1. Choose regions R, C and compute the archimedean functional

$$\mathcal{A}(R, C) = \frac{1}{2}\mathcal{E}(\mu_R) + \frac{1}{2}\mathcal{E}(\nu_C) + \frac{3}{4} + \int_0^1 Q_{\mu_R}(t)Q_{\nu_C}(t) dt,$$

with the scale normalization determined by $\text{area}(R) = \text{area}(C)$.

2. Compute or bound the exact prime transport functional

$$\mathcal{P}(R, C; A, B) = \sum_p \log p \inf_{\pi} \int w_p(r, \pi(r)) dr.$$

3. Any counterexample must satisfy $\mathcal{A}(R, C) \geq \mathcal{P}(R, C; A, B)$ for every admissible pair of regions. A single strict reverse inequality proves a contradiction.

Rectangles and origin triangles are analytically soluble test cases. The next step is not merely to scan aspect ratios: it is to optimize jointly over region shape and valuation configurations allowed by the two-three unit corollary and the primitive-gcd lemma of Section 5. The latter constraints are discrete and may prevent the valuation directions that make the transport cost artificially small.

59 Confluent and integer-valued determinant extensions

Ordinary evaluations use one row per lattice point. A potentially stronger construction uses integer jets. View (222) as evaluating the monomials $X^n Y^k$ at the integer points

$$(X, Y) = (2^a 3^b, A^a B^b).$$

The Euler operators

$$D_X = X \frac{\partial}{\partial X}, \quad D_Y = Y \frac{\partial}{\partial Y}$$

act by

$$D_X^r D_Y^s (X^n Y^k) = n^r k^s X^n Y^k,$$

which remains integral at the lattice points. Hasse derivatives give the alternative integer values

$$\binom{n}{r} \binom{k}{s} X^{n-r} Y^{k-s}.$$

Thus one can build confluent evaluation determinants without introducing the irrational factors $n + k\beta$ that ordinary derivatives along the real curve would produce.

The unresolved issues are nonvanishing and optimal scaling. A useful target is an integer-jet determinant whose archimedean limit behaves like a confluent HCIZ determinant while its p -adic assignment problem acquires an additional N^2 contribution. This is more promising than generic factorial normalization: the elementary divisibility obtained from binomial bases alone is only of order $N^{3/2} \log N$ for balanced two-dimensional grids.

60 Conditional resolutions and exact reformulations

60.1 Four Exponentials and Schanuel

The Four Exponentials Conjecture immediately implies Conjecture 2.1. Schanuel’s conjecture implies Four Exponentials and therefore also suffices. The present problem is thus a particularly concrete integral special case of a central algebraic-independence conjecture.

60.2 Logarithm determinant formulation

With $A = 2^x$ and $B = 3^x$, the condition is

$$\det \begin{pmatrix} \log 2 & \log A \\ \log 3 & \log B \end{pmatrix} = 0.$$

The conjecture says that if A, B are positive integers and this determinant vanishes, then the two columns are related by a nonnegative integer factor. Baker’s theorem controls nonzero *linear* forms in logarithms, whereas this is a bilinear relation among four logarithms; that distinction is exactly where standard transcendence machinery stops.

60.3 Order-preserving monoid formulation

A counterexample would give an order-preserving monoid isomorphism

$$\langle 2, A \rangle_{\mathbb{N}_0^2} \longrightarrow \langle 3, B \rangle_{\mathbb{N}_0^2}, \quad 2^n A^k \mapsto 3^n B^k,$$

because both sides have logarithm $n + k\beta$ up to a fixed positive scale. The exact primitive-generator theorem says this monoid is the entire domain on which $m \mapsto m^{\log_2 3}$ is integer-valued. A purely arithmetic classification of such order-preserving power isomorphisms would prove the conjecture without explicit transcendence estimates.

61 Formalization roadmap

A substantial portion of the new work is suitable for formalization in Lean, with the deep analytic inputs isolated behind explicit interfaces.

The first three items of the source roadmap are already discharged in the repository and are therefore not repeated as targets: the range, rationality, additive-monoid and local-finiteness lemmas are `IntegerExponent` and `LeastSolution`; the affine rank-two proposition is `MultiplicativeRank` and `AlgebraicExponentLocus`; and the primitive-gcd and denominator-elimination lemmas, together with the exact monoid identity they yield, are `PrimitiveOddCore`, `ThreeDenominatorNormalization`, `RationalPowerIndex` and `PrimitiveGenerator`. What remains is the following.

1. Define finite matrices (222). The tropical valuation lemma follows directly from the Leibniz formula and valuation of a sum. The square-grid assignment value (225) is a finite rearrangement theorem.
2. Treat strict total positivity, HCIZ, logarithmic-energy convergence, and asymptotic transport as separate analytic interfaces. Finite versions of all determinant inequalities can still be machine checked once supplied with rational interval certificates.
3. For computational searches over region shapes or rational bases, emit proof certificates consisting of finite node lists, exact valuation assignment dual variables, and rational upper bounds for the real determinant.

This separation would let the ProveIt development verify every new arithmetic step even before the transcendence and asymptotic libraries are available.

62 Prioritized research directions

1. **Joint valuation optimization.** Replace the componentwise assignment lower bound by the exact transport cost for each prime, and optimize over valuation vectors satisfying the primitive-gcd and base-prime-unit conditions. This is the most direct way to seek the missing N^2 constant.
2. **Rational-basis denominator network.** Apply Theorem 54.1 to a finite set of continued-fraction bases. Translate the resulting inequalities into lower bounds on denominators of B^q/A^p . Try to show that unique factorization makes the bounds mutually incompatible.

3. **Variational region search.** Numerically optimize convex and nonconvex lattice regions using the exact archimedean energy and assignment functionals. Any promising negative gap should then be converted into a rigorous interval proof.
4. **Integer jets and confluent total positivity.** Establish a nonvanishing theorem for Euler/Hasse-jet determinants and determine whether their prime transport grows faster than their archimedean capacity.
5. **Exploit minimal-generator valuations.** The conditions $\min(v_2(A), v_3(B)) = 0$ and $g_0 = 1$ have not yet been fully inserted into the variational optimization. They should be imposed from the beginning rather than used only as post-processing constraints.
6. **Effective finite certificates.** For candidate ranges of A , combine rigorous interval arithmetic with determinant and gcd constraints. Although finite verification cannot settle the conjecture alone, it can reveal which valuation patterns come closest to saturating the theoretical bounds.

63 A proof-producing finite verification

The structural and determinant arguments above are uniform, but they leave an infinite range of possible values of $A = 2^\beta$. A complementary finite calculation can remove an initial segment without assuming any unproved transcendence statement. The following computation is deliberately based on outward-rounded intervals and an explicit positive-series remainder, rather than on a test such as “the floating-point value is not visually close to an integer.”

Theorem 63.1 (Certified finite range). *For every integer A with*

$$2 \leq A \leq 100,000,$$

the number $A^{\log_2 3}$ is an integer if and only if A is a power of 2. Consequently, a nonintegral Alaoglu–Erdős counterexample must satisfy

$$2^\beta > 100,000, \quad \beta > \log_2(100,000) \approx 16.609640474437.$$

Certificate method. For fixed A , put

$$F_A(b) = (\log b)(\log 2) - (\log A)(\log 3), \quad b > 0.$$

The function F_A is strictly increasing, and $A^{\log_2 3} = b$ is equivalent to $F_A(b) = 0$. Every logarithm is enclosed without calling a transcendental floating-point routine in the decisive comparison. For $0 \leq z \leq 1/3$ we use

$$\log \frac{1+z}{1-z} = 2 \sum_{j=0}^{M-1} \frac{z^{2j+1}}{2j+1} + R_M(z), \quad 0 < R_M(z) < \frac{2z^{2M+1}}{(2M+1)(1-z^2)}.$$

For an integer $n \geq 1$, choose $k = \lfloor \log_2 n \rfloor$ and write

$$\log n = k \log 2 + \log \frac{1+z_n}{1-z_n}, \quad z_n = \frac{n-2^k}{n+2^k} \in [0, 1/3).$$

Thus all terms are positive rational operations, and directed rounding gives certified lower and upper endpoints. The run used 70 decimal digits and $M = 28$ terms. A binary64 value was used only as

an initial guess for the integer immediately below the root; the program then moved monotonically until it had certified, for consecutive integers $b, b + 1$,

$$F_A(b) < 0 < F_A(b + 1).$$

Hence no integer can be the root. Powers of 2 are separated at the start, and for $A = 2^m$ the corresponding value is exactly 3^m .

The completed run checked 99,983 nonpowers of 2. The largest number of integer steps needed to repair an initial binary64 guess was 0. The smallest absolute interval endpoint encountered was 6.263278510736485942331545732E-13, far outside the interval uncertainty. Runtime on the working container was approximately 15.238 seconds. The exact verifier used for the run is reproduced in Appendix Q; its SHA-256 digest is

acc099fb776f830b243a2421bc7ce27d5bedc63a38c3a7b190accb3c621cbadd.

□

Remark 63.2. The numerical bound is not presented as evidence that near misses cannot occur later. Its role is narrower and rigorous: it turns all A in the displayed finite interval into closed cases, while the determinant inequalities address whole infinite parameter ranges and arithmetic configurations at once.

64 Conclusion

The continuation does not close the final four-exponentials gap, but it substantially narrows and clarifies the problem.

First, the hypothetical exceptional set is completely rigid: one least counterexample β would generate every solution, and only solutions, as $n + k\beta$ with $n, k \geq 0$. Second, that exact grid supports determinants that are simultaneously strictly totally positive over the reals and highly divisible at finite primes. The square-grid divisor and the triangular inequality quantify both sides at the correct N^2 scale. Third, rational basis reduction is no longer a vague temptation: the adelic inequality shows exactly how denominators pay for every smaller logarithm. Finally, the remaining obstacle has a precise form—one needs an additional N^2 -scale arithmetic gain, not merely more lattice points or a sharper polynomial error term.

The most promising immediate attack is therefore a joint one: optimize the triangular or variational determinant using the full prime valuation vectors, while applying several rational bases to force incompatible denominator and gcd requirements. That program is concrete, computationally testable, and amenable to partial formal verification.

65 Third continuation: matrix coefficients, sparse interpolation, and boundary defects

Reports 1–2 in [Article/new-reports/](#) are independent continuations of the same unified base and therefore share most of their inherited body with this document. They are reconciled here by mathematical content rather than by concatenation: stale module statuses, the old candidate-count

off-by-one, and the incorrect first-layer valuation equality are not reintroduced. Their genuinely new core is a matrix-coefficient reformulation, an exact sparse-interpolation threshold, a Newton-polygon degree calibration, and a boundary-defect identity for the extremal interpolant. The matrix coefficient, formal prime-symbol, and sparse-interpolation components now have dedicated kernel-checked Lean modules.

Status: The full conjecture remains open. The restricted matrix-coefficient equivalence, formal prime-symbol zero classification, Structural Rank equivalence, finite sparse power-curve lemmas, and boundary-dilation algebra below are [\[kernel-verified Lean\]](#). Newton-polygon geometry, the packaged BK-degree threshold, the counting reinterpretations, and the complex monodromy consequence remain conventional mathematics or consequences of named classical inputs.

65.1 The restricted matrix-coefficient principle

For positive integers M, A , put

$$\mathbf{L}(M, A) = \begin{pmatrix} \log 2 & \log M \\ \log 3 & \log A \end{pmatrix}.$$

The determinant vanishes exactly when there is a real β with

$$M = 2^\beta, \quad A = 3^\beta.$$

In that case

$$\mathbf{L}(M, A) = \begin{pmatrix} \log 2 \\ \log 3 \end{pmatrix} \begin{pmatrix} 1 & \beta \end{pmatrix},$$

and for rational vectors $w = (r, s)^T$ and $v = (u, z)^T$ one has the exact factorization

$$w^T \mathbf{L}(M, A) v = (r \log 2 + s \log 3)(u + z\beta). \quad (239)$$

Define $\text{MCC}_{2,3}^{\mathbb{Z}}$ to assert that every singular matrix of this restricted form has nonzero rational w, v with $w^T \mathbf{L}(M, A) v = 0$.

Theorem 65.1 (Restricted Matrix Coefficient equivalence). *The Alaoglu–Erdős conjecture is equivalent to $\text{MCC}_{2,3}^{\mathbb{Z}}$.*

Proof. If the conjecture holds, singularity supplies a common exponent β , which must be an integer m ; the rational right vector $(-m, 1)^T$ lies in the kernel. Conversely, suppose β is a nonintegral common exponent. Integrality of 2^β implies that β is irrational. In (239) the left factor cannot vanish for nonzero rational w , by multiplicative independence of 2 and 3, and the right factor cannot vanish for nonzero rational v , by irrationality of β . Thus the required zero matrix coefficient does not exist. $\square$

The new module `RestrictedMatrixCoefficient` formalizes the concrete 2×2 matrix, its determinant, the common-exponent equivalence, the rank-one identity `restrictedLogMatrix_eq_rankOne`, the coefficient factorization and anisotropy, and the exact endpoint

`alaogluErdosConjecture_iff_restrictedMatrixCoefficientPrinciple.`

This is [\[kernel-verified Lean\]](#). It is a strict specialization of the Matrix Coefficient Conjecture of Dasgupta–Kakde [37]; in the real 2×2 case their fewnomial condition implies the desired matrix coefficient, while constructing that fewnomial from singularity remains the open direction.

65.2 The formal prime-symbol determinant

Let $a_p = v_p(M)$ and $b_p = v_p(A)$, and introduce one formal variable X_p for each prime in $\{2, 3\} \cup \text{supp}(MA)$. The formal determinant is

$$\Phi_{M,A}(\mathbf{X}) = X_2 \sum_p b_p X_p - X_3 \sum_p a_p X_p. \quad (240)$$

Substitution $X_p = \log p$ gives $\det L(M, A)$. Coefficient comparison gives

$$\Phi_{M,A} \equiv 0 \iff \exists m \in \mathbb{N}_0, M = 2^m \text{ and } A = 3^m. \quad (241)$$

Indeed, all b_p away from 3 and all a_p away from 2 must vanish, while the $X_2 X_3$ coefficient forces $b_3 = a_2$. Hence the conjecture is equivalently the assertion that specialization of this signed sparse quadratic at the prime-logarithm point cannot vanish unless the polynomial is formally zero.

For a hypothetical counterexample $\Phi_{M,A}$ is nonzero and irreducible over $\mathbb{Q}$. A factorization of a homogeneous quadratic would yield two rational linear factors; one of them would vanish after $X_p = \log p$, contradicting rational linear independence of logarithms of distinct primes. The new `FormalPrimeDeterminant` module defines the genuine `MvPolynomial formalPrimeDeterminant`, proves its specialization `formalPrimeDeterminant_evalâĈĈ_log`, the unconditional classification `formalPrimeDeterminant_eq_zero_iff`, and the exact equivalence `alaogluErdosConjecture_iff_sparsePrimeLogStructuralRankPrinciple`. These claims are [\[kernel-verified Lean\]](#). The homogeneous-factor irreducibility bookkeeping is intentionally kept separate and remains paper-level.

65.3 Exact sparse interpolation on an irrational power curve

For irrational β , a monomial $X^r Y^s$ restricted to $(X, Y) = (e^t, e^{\beta t})$ becomes

$$e^{(r+s\beta)t}.$$

Distinct pairs $(r, s) \in \mathbb{N}_0^2$ give distinct frequencies. The classical Chebyshev zero theorem for real exponential polynomials, already formalized in `ExponentialPolynomialZeros`, therefore gives the following sharp lower bound.

Theorem 65.2 (Sparse power-curve zero bound). *An m -term nonzero polynomial restricted to $Y = X^\beta$, with β irrational, has at most $m - 1$ distinct positive zeros. Equivalently, vanishing at R distinct positive curve points requires at least $R + 1$ nonzero monomials.*

The new `SparsePowerCurve` module proves the arbitrary-order indexed form `sparsePowerCurve_coeff_eq_zero_o`, its finite-zero-set form `card_le_of_sparsePowerCurve_eq_zero`, and the nonsingularity of the square evaluation matrix. It also proves the integral form `det_natPowerCurveEvaluationMatrix_ne_zero`: if $y_i = x_i^\beta$ are natural numbers, then

$$\det(x_i^{r_j} y_i^{s_j}) \neq 0$$

for distinct monomial pairs and ordered positive x_i . Thus the determinant is a nonzero integer and has absolute value at least one. These statements are [\[kernel-verified Lean\]](#).

The matching upper construction is the node polynomial

$$P_0(T) = \prod_{i=1}^R (T - x_i).$$

Because all x_i are positive, every elementary symmetric coefficient is nonzero, so P_0 has exactly $R + 1$ monomials. This coefficient assertion and the vanishing at every node are also formalized as `positiveNodeInterpolant_support_card` and `positiveNodeInterpolant_eval`. Consequently the following packaged minimum is a paper-level corollary whose two substantive halves are kernel checked:

$$\min\{\#\text{supp } P : 0 \neq P \in \mathbb{C}[X, Y], \ P(x_i, x_i^\beta) = 0 \ (1 \leq i \leq R)\} = R + 1. \quad (242)$$

Here the paper-level minimum allows complex coefficients; the kernel-checked zero-bound and matching positive-node construction are stated over the real core needed in the application.

In the restricted integer-output setting this gives an exact dichotomy. If the common exponent is an integer m , the minimum support on any nonempty family is 2, attained by $Y - X^m$; a single nonzero monomial cannot vanish at a positive point. If the exponent is nonintegral, it is irrational and the minimum on R distinct points is $R + 1$ by (242). The irrational branch is the kernel-checked sparse theorem; the integral branch is the displayed elementary two-term witness and positivity argument.

For a counterexample (M, A, β) and a finite $\Omega \subset \mathbb{N}_0^2$, the points

$$(2^a 3^b, M^a A^b) = (e^{z_{a,b}}, e^{\beta z_{a,b}}), \quad z_{a,b} = a \log 2 + b \log 3,$$

have distinct parameters. Hence $R = \#\Omega$ gives the exact minimum $R + 1$. On the $2N \times 2N$ box, any polynomial with at most $4N^2$ monomials would therefore save one term over trivial interpolation and rule out a counterexample. Dasgupta–Kakde ask for the stronger support bound $< N^2$ on the same $4N^2$ points; for Alaoglu–Erdős, one saved monomial already suffices.

65.4 BK degree and the exact factor-four threshold

For $0 \neq P \in \mathbb{Q}[X, Y]$, first divide out the greatest common monomial of its terms, as in the Dasgupta–Kakde translation-invariant normalization. Let P^{red} denote the resulting polynomial, let $\Delta(P)$ be its Newton polygon, and set

$$\text{bkd}(P) = \max\{\deg P^{\text{red}}, 2 \text{Area } \Delta(P)\}.$$

Pick’s theorem gives

$$\#\text{supp } P \leq \text{bkd}(P) + 2, \quad (243)$$

with $+1$ when the Newton polygon is one-dimensional. For irrational β and $R \geq 2$ distinct positive rational nodes whose ordinates x_i^β are rational, let $L \in \mathbb{Q}[X]$ be the Lagrange interpolation polynomial of degree at most $R - 1$ with $L(x_i) = x_i^\beta$. Combining (242) with the thin interpolant $Y - L(X)$ forces all $R + 1$ available coefficients to occur and yields the exact paper-level threshold

$$\min\{\text{bkd}(P) : 0 \neq P \in \mathbb{Q}[X, Y], \ P(x_i, x_i^\beta) = 0\} = R - 1. \quad (244)$$

For the $K \times K$ semigroup grid with $K \geq 2$ this is $K^2 - 1$; on the Dasgupta–Kakde $2N$ grid it is $4N^2 - 1$. Their target $< N^2$ must therefore cross a genuine asymptotic factor-four gap. No Pick-theorem or normalized lattice-area library is currently present in the pinned Mathlib, so (243)–(244) remain [\[paper proof in this report\]](#).

65.5 Boundary defects of the extremal interpolant

Assume $K \geq 2$ and a counterexample normalization $M = 2^\beta$, $A = 3^\beta$ with β irrational. Let $Q_K \in \mathbb{Q}[X]$ be the degree- $K^2 - 1$ polynomial interpolating the second coordinate on the $K \times K$ grid, and let

$$\Pi_{r,s}(X) = \prod_{\substack{0 \leq i < r \\ 0 \leq j < s}} (X - 2^i 3^j).$$

The interior grid is stable under the two dilations. Therefore

$$Q_K(2X) - MQ_K(X) = \Pi_{K-1,K}(X)U_K(X), \quad (245)$$

$$Q_K(3X) - AQ_K(X) = \Pi_{K,K-1}(X)V_K(X), \quad (246)$$

where U_K, V_K are unique of degree $K - 1$; their leading coefficients are respectively $q_K(2^{K^2-1} - M)$ and $q_K(3^{K^2-1} - A)$ when q_K is the leading coefficient of Q_K . The irrationality of β makes both displayed factors nonzero. More generally, the same degree conclusion requires the explicit noncancellation hypotheses $2^{K^2-1} \neq M$ and $3^{K^2-1} \neq A$.

The dilation operators $D_2f = f(2X) - Mf(X)$ and $D_3f = f(3X) - Af(X)$ commute. Substituting (245)–(246) and cancelling the common interior factor $\Pi_{K-1,K-1}$ produces an identity of degree at most $2K - 2$ involving the top and right boundary pieces together with their bottom and left dilation preimages. This boundary compression is a concrete finite target: it turns K^2 interpolation constraints into a low-degree compatibility equation. It does not by itself produce the missing fewnomial, and no such claim is made.

The module `BoundaryDilationDefect` proves the residual polynomial algebra over a general field and the commutator identity over an infinite field. Its theorem `existsUnique_squareGridResiduals` gives both unique residuals, their exact degree and leading coefficients; `boundaryDilation_commutator` proves the displayed commuting boundary identity; and `squareBoundaryExpressions_natDegree_le` proves the $2K - 2$ bounds. The interpolation values, node injectivity, and noncancellation of the top coefficient remain visible hypotheses. These statements are [\[kernel-verified Lean\]](#).

65.6 Counting thresholds and rational monodromy

A counterexample gives the rational points

$$P_{n,k} = (2^n M^k, 3^n A^k) = (2^{n+k\beta}, 3^{n+k\beta}), \quad n, k \in \mathbb{N}_0,$$

on the fixed curve $Y = X^\vartheta$. Since the second coordinate controls the height, a triangular lattice count gives

$$\#\{(n, k) \in \mathbb{N}_0^2 : n + k\beta \leq \log_3 H\} = \frac{(\log H)^2}{2\beta(\log 3)^2} + O_\beta(\log H).$$

Thus any uniform $o((\log H)^2)$ upper bound for rational points on the unbounded power curve would prove the conjecture. The synchronized compact orbit from `CompactOrbit` gives a sharper

arithmetic target: its first k points have height at most A^k , so an $o(\log H)$ bound for points retaining the common floor level would also suffice. These are height reinterpretations of the accepted exact and quadratic conditional counts, not a new point-counting theorem.

Finally, classical six exponentials forces a complex consequence which the current real determinant formalization does not cover. If β were a nonintegral solution and $q \in \mathbb{Q}^\times$, then

$$e^{2\pi i q \beta} \text{ is transcendental,} \tag{247}$$

and so are its sine and cosine. Apply six exponentials to the independent rows $1, \beta$ and the independent columns $\log 2, \log 3, 2\pi i q$; five of the six exponentials are algebraic, forcing the sixth to be transcendental. This is a valid consequence of the named classical theorem, not a new kernel theorem.

65.7 The reported 2^{29} scan and its trust boundary

The first and fifth new reports record an MPFR directed-rounding extension through 2^{29} , including chunk tables, SHA-256 identifiers, extremal records, and high-precision replays; report 5 also embeds the literal C source and a full manifest in its TeX. The logical consequence would be that a nonintegral solution has $x > 29$. However, the merged **new-reports/** directories contain no standalone scanner, canonical chunk logs, or replay artifacts named by those hashes, and the long run was not independently repeated as part of this merge. It is therefore retained as a *reported [software computation]* result, while the repository-reproducible software frontier in Section 40 remains 2^{27} .

65.8 What changed, and what did not

The continuation produces exact finite reformulations, but no unconditional proof of the conjecture. Its most useful hierarchy is now:

1. the conjecture is exactly the restricted matrix-coefficient principle (**[kernel-verified Lean]**);
2. the formal prime-symbol determinant specializes to the logarithmic determinant, and its sparse Structural Rank principle is again exactly the conjecture (**[kernel-verified Lean]**);
3. a counterexample has maximal sparse interpolation complexity $R + 1$, with the lower bound, matching positive interpolant, and evaluation determinant kernel checked;
4. the corresponding BK-degree threshold is exactly $R - 1$ (**[paper proof in this report]**), exposing a factor-four gap to the current Matrix Coefficient auxiliary-polynomial target;
5. dilation defects compress the extremal grid interpolant to boundary polynomials of degree $K - 1$ (**[kernel-verified Lean]**), but an arithmetic construction saving one monomial remains open;
6. generic polylogarithmic point counting is insufficient unless its exponent is below two globally, or below one on the synchronized compact orbit.

The Massold v1 parameter substitution audited in report 2 [38] does not cover the 2×2 case: its displayed choice gives $t = 1$, $\mu = \nu = 1$, making the required strict inequality $1/2 > 1$. This is a narrow provenance warning about that version, not a claim about every possible correction.

66 Further continuations: support spectra and arithmetic determinant audits

Reports 3–5 are likewise independent continuations of that common base. This section records their new mathematics after reconciling overlaps rather than treating the reports as a sequence. Three themes survive audit: the exact support spectrum of a finite orbit, integer-valued-polynomial divisors and primitive Smith normalization, and an adelic audit of exact-grid determinants including a signed-progression no-go theorem. We do not import a stated “quadratic divisor criterion” from report 5: an inequality containing an unspecified $o(N^2)$ term is not a sufficient eventual strict inequality. Likewise, the proposed cusp lattice realization is kept conditional on the uniform logarithmic-energy convergence that its current proof does not establish.

66.1 The exact orbit-collision support spectrum

For a finite nonempty $\Omega \subset \mathbb{Z}^2$, define

$$T_x(\Omega) = \{i + jx : (i, j) \in \Omega\}, \quad D_x(\Omega) = \#T_x(\Omega),$$

and let $\mu_x(\Omega)$ be the least number of monomials in a nonzero Laurent polynomial in $\mathbb{C}[X^{\pm 1}, Y^{\pm 1}]$ vanishing at all points

$$X_x(\Omega) = \{2^{i+jx}, 3^{i+jx} : (i, j) \in \Omega\}.$$

Multiplication by a monomial shows that the same minimum is obtained with ordinary polynomials. Restriction to the orbit turns every monomial into an exponential with frequency $a \log 2 + b \log 3$; these frequencies are distinct by unique factorization. Consequently the exponential-polynomial zero bound and the positive-node polynomial give the exact formula

$$\boxed{\mu_x(\Omega) = D_x(\Omega) + 1.} \tag{248}$$

More generally, every evaluation matrix over $\mathbb{C}$ has rank $\min(D_x(\Omega), \#S)$ for any finite monomial support S . Repeated zeros counted with multiplicity obey the same rule: prescribing total multiplicity R costs at least $R + 1$ monomials.

For $L, K \geq 1$, on the rectangle $\Omega_{L,K} = \{0, \dots, L-1\} \times \{0, \dots, K-1\}$ the collision count is explicit. If $x \notin \mathbb{Q}$, then $D_x(L, K) = LK$. If $x = r/s \in \mathbb{Q}_{\geq 0}$ is in lowest terms, then

$$D_{r/s}(L, K) = LK - (L-r)_+(K-s)_+, \quad \mu_{r/s}(L, K) = D_{r/s}(L, K) + 1. \tag{249}$$

Indeed two vertices collide exactly along the primitive direction $(r, -s)$; the collision classes are paths, and vertices minus edges gives (249). Thus rationality is equivalent to saving one monomial on some square grid, while an irrational exponent has $\mu_x(L, L) = L^2 + 1$ for every L . For integer outputs, a complex dependence may be replaced by a rational one and denominators cleared, giving an integer certificate. This is the exact finite-support version of the restricted Matrix Coefficient target.

The module `RationalOrbitGrid` kernel-checks the exact rectangle formula, including the zero-numerator edge case, the irrational full-cardinality statement, the irrational square-grid sparse lower bound, and the canonical one-monomial-saving rational-orbit certificate once the grid exceeds the reduced numerator and denominator. Together with `SparsePowerCurve`, this certifies the real, nonnegative-exponent core of the spectrum above. The arbitrary finite Laurent support/rank formulation and the confluent multiplicity packaging remain paper-level. [\[kernel-verified Lean\]](#)

66.2 Coefficient height and the generic-modulus no-go

The exact support count does not control coefficient height. For $M \geq 1$, under an irrational exponent x with $A = 2^x \in \mathbb{Z}_{>0}$, the canonical annihilator of the M^2 first-coordinate nodes is

$$Q_M(T) = \prod_{0 \leq i, j < M} (T - 2^i A^j) \in \mathbb{Z}[T].$$

It has exactly $M^2 + 1$ nonzero coefficients. Write $H(Q) = \max_k |q_k|$ for $Q(T) = \sum_k q_k T^k$. Writing $C_M = (2A)^{M^2(M-1)/2}$, positivity of the roots gives the exact finite bounds

$$C_M \leq \|Q_M\|_1 \leq 2^{M^2} C_M, \quad \frac{C_M}{M^2 + 1} \leq H(Q_M) \leq 2^{M^2} C_M, \quad (250)$$

and hence

$$\log H(Q_M) = \frac{(1+x) \log 2}{2} M^2(M-1) + O(M^2 + \log M).$$

Thus the canonical support is sharp but carries an $\exp(\Theta_x(M^3))$ archimedean cost. Adding confluent multiplicities does not lower the support cost: the exponential Chebyshev bound charges one extra monomial for each unit of prescribed total multiplicity.

The corresponding adelic criterion is elementary and exact. If $P \in \mathbb{Z}[X, Y]$ and, at every integer grid point z , a finite set of primes S satisfies

$$|P(z)|_\infty \prod_{p \in S} |P(z)|_p < 1,$$

then the product formula forces $P(z) = 0$. Equivalently, it is enough to find a divisor $q_z \mid P(z)$ with $|P(z)| < q_z$ at every row. This identifies the right target but does not construct the coefficients.

A cardinality-only modular pigeonhole does not suffice. Let $E \in \mathbb{Z}^{R \times K}$ with $1 \leq K \leq R$, let $Q \geq 2$ and $H \geq 0$, and assume $|E_{rj}| \leq M_r$ with $M_r \geq 1$. Using coefficient vectors in $\{0, \dots, H\}^K$ and one modulus Q requires $(H+1)^K > Q^R$, which forces $H \geq Q$; the generic bound $|(Ed)_r| \leq KHM_r$ can then never prove $|(Ed)_r| < Q$. Even row-dependent moduli do not repair this architecture: for $H \geq 1$ and $q_r \geq 2$, the two sufficient conditions

$$(H+1)^K > \prod_{r=1}^R q_r, \quad q_r > KHM_r \quad (1 \leq r \leq R)$$

are incompatible. Finally, if one uses the full adelic product—or a finite prime set containing every prime divisor of the row content g_r —then dividing that row by g_r changes neither the archimedean- p -adic product nor the meaningful cancellation: the product formula for g_r exactly cancels the apparent common divisibility. The useful arithmetic object is therefore the Smith profile of the row-primitive evaluation matrix, not a raw common modulus.

These coefficient-height, adelic, pigeonhole, and row-content statements are paper-level. The same support proof is base-independent for every multiplicatively independent positive base tuple: distinct monomials restrict to exponentials with distinct real frequencies.

66.3 Universal factorial content and primitive interpolation

For $r \geq 1$, write

$$\mathfrak{F}_r = \prod_{j=0}^{r-1} j!.$$

For arbitrary integers $m_0, \dots, m_r, y_0, \dots, y_r$, the interpolation determinant with columns $1, m, \dots, m^{r-1}, y$ is divisible by $\mathfrak{F}_r$. Replacing the power columns by $\binom{m}{0}, \dots, \binom{m}{r-1}$ proves the claim by a triangular Stirling change of basis, and the nodes $0, \dots, r$ with ordinate vector $(0, \dots, 0, 1)$ show that the universal divisor is best possible. Applied to candidate ordinates, this strengthens the arbitrary-node repulsion constant by the factor $\mathfrak{F}_r^{1/\binom{r+1}{2}}$ but leaves its exponent unchanged.

The focused module **SuperfactorialDeterminant** kernel-checks the nonnegative-integer-node instance of this universal divisor, its exact consecutive-node sharpness witness, the resulting lower bound for a nonzero candidate interpolation determinant, and the strengthened divided-difference repulsion and packing inequalities. As in **ArbitraryNodeRepulsion**, the final two endpoints retain **RpowDividedDifferenceMeanValue** as an explicit hypothesis; the superfactorial divisibility itself is unconditional in that formalized scope. The arbitrary signed-integer node formulation above remains a paper consequence. [\[kernel-verified Lean\]](#)

The bivariate version is equally exact. If $\mathcal{S} \subset \mathbb{N}_0^2$ is a finite coordinate lower set, then for exactly $|\mathcal{S}|$ integer points (X_i, Y_i) ,

$$\prod_{(a,b) \in \mathcal{S}} a!b! \mid \det(X_i^a Y_i^b)_{i,(a,b) \in \mathcal{S}}. \quad (251)$$

The proof uses the integer-valued basis $\binom{X}{a} \binom{Y}{b}$. On a weighted triangular lower set of cardinality N , the logarithm of this divisor is $O(N^{3/2} \log N) = o(N^2)$. It is a genuine source of primes absent from the matrix entries, but it cannot by itself repair a positive quadratic-scale determinant screen.

There is a sharper invariant for the unique $(M+1)$ -term interpolating relation. Let $M \geq 2$ and let $u_1 < \dots < u_M$ be distinct integers, with $v_1, \dots, v_M \in \mathbb{Z}$. For the $M \times (M+1)$ matrix whose i th row is $(v_i, 1, u_i, \dots, u_i^{M-1})$, let Δ_M be the gcd of all signed maximal minors and let

$$G_M = \gcd_i \prod_{a < b, a, b \neq i} (u_b - u_a), \quad V_M = \prod_{a < b} (u_b - u_a), \quad R_i = \prod_{j \neq i} |u_i - u_j|.$$

Then a primitive cofactor generator, determined only up to global sign, has Y -coefficient satisfying $|c_Y| = V_M / \Delta_M$, and

$$\mathfrak{F}_{M-1} \mid G_M \mid \Delta_M, \quad G_M = \frac{V_M}{\text{lcm}_i R_i}. \quad (252)$$

Moreover G_M is the greatest divisor forced uniformly by the nodes as the integer ordinate vector varies. The remaining quotient Δ_M / G_M is therefore the precise output-dependent Smith mass that an arithmetic proof must enlarge. These Smith statements and their $2^i 3^j$ -grid asymptotics remain paper-level.

66.4 The adelic–HCIZ accounting identity

Let $N \geq 1$. For an exact-grid matrix $E_{ij} = \exp(u_i v_j) \in \mathbb{N}$ with strictly increasing $u_1 < \dots < u_N$ and $v_1 < \dots < v_N$, let $D_N > 0$ be its determinant. Put

$$C_N^- = \sum_i u_i v_{N+1-i}, \quad M_N = \frac{(\sum_i u_i)(\sum_j v_j)}{N}, \quad C_N^+ = \sum_i u_i v_i.$$

For each prime let τ_p be the minimum valuation of a Leibniz term, and set $T_N = \sum_p \tau_p \log p$ and $Q_N = \prod_p p^{\tau_p}$. The HCIZ identity writes

$$\log D_N = V_N + \log I_N, \quad V_N = \log \Delta(u) + \log \Delta(v) - \log \prod_{j=1}^{N-1} j!.$$

Here I_N is the Harish–Chandra–Itzykson–Zuber integral $\int_{U(N)} \exp(\text{tr}(\text{diag}(u)U \text{diag}(v)U^*)) dU$ for normalized Haar probability measure. Rearrangement, averaging, Jensen, and von Neumann’s trace inequality give the exact sandwich

$$T_N \leq C_N^- \leq M_N \leq \log I_N \leq C_N^+. \quad (253)$$

Defining

$$\Xi_N = \log(D_N/Q_N), \quad S_N = V_N + M_N - C_N^-, \quad H_N = \log I_N - M_N, \quad A_N = C_N^- - T_N$$

therefore gives

$$\boxed{\Xi_N = S_N + H_N + A_N}, \quad H_N, A_N, \Xi_N \geq 0. \quad (254)$$

This separates geometric screen, random-matrix entropy, and incompatible primewise assignments instead of hiding them in one determinant estimate.

For the inherited triangular spectra, the best possible joint screen has normalized limit

$$f_{\text{joint}}(c) = \frac{1}{2} \log c - 1 + \left(\frac{4}{9} - \frac{\pi}{8}\right) c, \quad c = 2\sqrt{\beta \log 2 \log 3}.$$

Its unique zero corresponds to $\beta = 6.9269442924 \dots$. Hence throughout the kernel-verified range $\beta \geq 12$, no combination of an exact HCIZ evaluation and termwise prime valuations can close the inherited triangular determinant family. This is a no-go theorem for that geometry, not for all determinant methods.

Report 5 proposes power-cusp spectra whose formal continuum screen tends to $-\infty$. The continuum formula itself is valid under convergence of logarithmic energies, but weak lattice convergence does not supply that hypothesis near the diagonal. A complete lattice realization needs an explicit spacing/truncation estimate. Even after such a repair, the separated-prime-support branch makes the adelic alignment defect grow linearly in the cusp parameter, so cusp concentration alone cannot be a universal proof.

66.5 Signed progressions and the denominator tax

Assume $N \geq 1$, $p, q, r, s, M, A \in \mathbb{N}_{>0}$, and the counterexample relations $M = 2^\beta$, $A = 3^\beta$ with β irrational. For $0 \leq i, j < N$, use the nonnegative exponent vectors

$$(a_i, b_i) = (iq, (N-1-i)p), \quad (n_j, k_j) = (jr, (N-1-j)s).$$

Set

$$X = 2^{qr}, \quad Y = 3^{pr}, \quad Z = M^{qs}, \quad W = A^{ps}, \quad P = XW, \quad Q = YZ.$$

Define the original integer matrix and its determinant by

$$E_{ij} = 2^{a_i n_j} 3^{b_i k_j} M^{a_i k_j} A^{b_i k_j}, \quad D_N = \det(E_{ij})_{0 \leq i, j < N},$$

and let Q_N^{trop} be the termwise tropical divisor of this matrix. Then

$$\log(P/Q) = (q \log 2 - p \log 3)(r - s\beta),$$

so $P \neq Q$: unique factorization makes the first factor nonzero, while irrationality makes the second nonzero. A direct reduction to a geometric Vandermonde gives

$$|D_N| = (PQ)^{\binom{N}{3}} \prod_{d=1}^{N-1} |P^d - Q^d|^{N-d}. \quad (255)$$

After writing $g = \gcd(P, Q)$, $\hat{P} = P/g$, $\hat{Q} = Q/g$, every row, column, and common gauge factor cancels from the tropical quotient:

$$\frac{|D_N|}{Q_N^{\text{trop}}} = \prod_{d=1}^{N-1} |\hat{P}^d - \hat{Q}^d|^{N-d}. \quad (256)$$

Let $B = \min(\hat{P}, \hat{Q})$ and $\eta = |\log(\hat{P}/\hat{Q})|$. With $G(N+1) = \prod_{d=1}^{N-1} d^{N-d}$, the exact expansion is

$$\log \frac{|D_N|}{Q_N^{\text{trop}}} = \binom{N+1}{3} \log B + \binom{N}{2} \log \eta + \log G(N+1) + R_N, \quad (257)$$

$$0 \leq R_N \leq \binom{N+1}{3} \eta. \quad (258)$$

The integer gap $|\hat{P} - \hat{Q}| \geq 1$ forces $\log B + \log \eta \geq -\eta$. Thus the product of two continued-fraction errors is not free: the denominator height repays the apparent quadratic smallness at cubic scale. This exact algebra rules out the naive double-continued-fraction progression, while leaving open more structured cyclotomic cancellation inside $\hat{P}^d - \hat{Q}^d$.

The module `SignedProgressionDeterminant` kernel-checks the geometric Vandermonde factorization, the primitive invariant quotient and its coprimality with both primitive bases, the exact finite logarithmic decomposition, the nonnegative remainder bound, and the integer-gap inequality. The original four-parameter matrix-to-tropical-quotient reduction, the continued-fraction asymptotics, and any general impossibility theorem for nontropical determinant cancellation remain paper-level. [\[kernel-verified Lean\]](#)

Status: The universal univariate superfactorial divisor and its strengthened repulsion/packing consequences are [\[kernel-verified Lean\]](#) in `SuperfactorialDeterminant`; the finite signed-progression and denominator-tax algebra is [\[kernel-verified Lean\]](#) in `SignedProgressionDeterminant`. The general support-spectrum, bivariate factorial/Smith, and HCIZ statements are paper proofs extracted from reports 3–5. For signed progressions, the original four-parameter/tropical identification and continued-fraction conclusion remain paper-level. No cusp lattice limit or full conjecture is asserted.

67 Sixth report: phase randomization and positive-kernel certificates

Report 6 is another independent continuation of the same common base. Its matrix-coefficient and sparse-grid arguments overlap Sections 65.1–65.3, so they are reconciled rather than duplicated. Its genuinely distinct contribution is a probability-theoretic control experiment that separates a generic random phase from the arithmetically distinguished phase zero. The experiment does not declare ϑ generic and does not prove the conjecture; it identifies one exact finite Fourier estimate that would suffice.

67.1 A uniform metric estimate and random-phase extinction

Fix $0 < a < b$ and, for $m \geq 2$ and $0 < \varepsilon \leq 1/4$, set

$$E_m(\varepsilon) = \{\alpha \in [a, b] : \|m^\alpha\|_{\mathbb{Z}} < \varepsilon\}.$$

The monotone substitution $y = m^\alpha$ and a harmonic-sum estimate give

$$\text{Leb}(E_m(\varepsilon)) \leq \frac{8}{3} \left(b + \frac{1 + \log 3}{\log 2} \right) \varepsilon. \quad (259)$$

The constant is independent of m . Consequently every nonnegative summable sequence $\sum_m \varepsilon_m < \infty$ satisfies

$$\|m^\alpha\|_{\mathbb{Z}} < \varepsilon_m \quad \text{for only finitely many } m$$

for almost every $\alpha > 0$, by the first Borel–Cantelli lemma. This is a metric theorem, not a specialization theorem for the named exponent ϑ .

For the synchronized model, fix $N \geq 1$ and put

$$D_N = \{a \in \mathbb{Z} : 2^N < a < 2^{N+1}\}, \quad Z_N(u; \eta) = \#\{a \in D_N : \|a^\vartheta - u\|_{\mathbb{Z}} < \eta\},$$

where u is a phase on $\mathbb{R}/\mathbb{Z}$. A uniform phase makes each radius- η arc have Haar measure 2η , so at the natural resolution $\eta = 3^{-N}$,

$$\mathbb{E}_u Z_N(u; 3^{-N}) = 2(2^N - 1)3^{-N} < 2(2/3)^N. \quad (260)$$

Summability and Borel–Cantelli show that almost every single phase has no sufficiently late near hits. No independence between blocks is used.

Phase zero is qualitatively different. Let

$$\mathcal{H}_N = \#\{a \in D_N : a^\vartheta \in \mathbb{Z}\}.$$

If the conjecture holds then $\mathcal{H}_N = 0$. Under failure, the least nonintegral generator β and **ExactCounting** give the exact nonintegral count

$$\mathcal{H}_N = \left\lfloor \frac{N+1}{\beta} \right\rfloor, \quad \frac{\mathcal{H}_N}{N} \longrightarrow \frac{1}{\beta}, \quad Z_N(0; 3^{-N}) \geq \mathcal{H}_N. \quad (261)$$

This is deliberately the nonintegral count: the full candidate count in the half-open dyadic block convention used elsewhere in the report has the additional integral point and equals $1 + \lfloor (N+1)/\beta \rfloor$.

67.2 The Fejér certificate

Write $e(t) = e^{2\pi it}$ and, for an integer $H \geq 2$, define

$$F_H(t) = \frac{1}{H} \left| \sum_{j=0}^{H-1} e(jt) \right|^2, \quad \Phi_H(t) = \frac{F_H(t)}{H}.$$

Then $\Phi_H \geq 0$, $\Phi_H(0) = 1$, and

$$\Phi_H(t) = \frac{1}{H} \sum_{|h| < H} \left(1 - \frac{|h|}{H} \right) e(ht).$$

With

$$S_N(h) = \sum_{a \in D_N} e(ha^\vartheta), \quad Q_N(H) = \sum_{a \in D_N} \Phi_H(a^\vartheta),$$

positivity and the finite Fourier expansion give the exact certificate

$$\mathcal{H}_N \leq Q_N(H), \tag{262}$$

$$Q_N(H) = \frac{|D_N|}{H} + \frac{2}{H} \sum_{h=1}^{H-1} \left(1 - \frac{h}{H}\right) \Re S_N(h). \tag{263}$$

Thus the conjecture follows if, along an unbounded sequence, one can choose $H = H_N^*$ with $|D_N|/H_N^* = o(N)$ and make the weighted Fourier term in (263) equal to $o(N)$. The natural bandwidth $H = 3^N$ has exponentially small random-phase mean, but it demands information at exponentially many frequencies; fixed-mode equidistribution from `BlockWeyl` is therefore far too coarse.

The same diagnosis appears in the exact second moment. If $d_{a,a'} = \left\| \{a^\vartheta\} - \{a'^\vartheta\} \right\|_{\mathbb{Z}}$ and $0 < \eta \leq 1/4$, then

$$\mathbb{E}_u Z_N(u; \eta)^2 = 2\eta |D_N| + 2 \sum_{a < a'} (2\eta - d_{a,a'})_+.$$

The second term is the microscopic pair-correlation energy that a counterexample must force to be large at phase zero.

67.3 Why naive random sparse coefficients do not bridge the gap

The integrality amplifier behind the sparse certificate is exact: if $D \in \mathbb{N}_{>0}$, $0 \neq P \in \mathbb{Q}[X, Y]$, $DP \in \mathbb{Z}[X, Y]$, and $|P(g)| < D^{-1}$ at every conditional integer-grid point g , then $P(g) = 0$ everywhere. Sections 65.3 and 66.1 then turn a subcritical support into rationality of the common exponent.

Unconditioned random signs do not naturally reach this threshold. For distinct monomials M_ν , real c_ν , and independent signs ϵ_ν , put $P_\epsilon = \sum_\nu \epsilon_\nu c_\nu M_\nu$. On every finite grid $G \subset \mathbb{R}_{>0}^2$,

$$\mathbb{E} \sum_{g \in G} |P_\epsilon(g)|^2 = \sum_\nu |c_\nu|^2 \sum_{g \in G} |M_\nu(g)|^2. \tag{264}$$

At a fixed g , if $\sigma_g^2 = \sum_\nu |c_\nu M_\nu(g)|^2$, the fourth-moment bound and Paley–Zygmund give

$$\Pr\{|P_\epsilon(g)| \geq \sigma_g/\sqrt{2}\} \geq \frac{1}{12}.$$

When $\sigma_g = 0$ this is immediate; otherwise it is the usual Paley–Zygmund application to $|P_\epsilon(g)|^2$. This gives an RMS-scale baseline at each fixed point. It does not exclude rare simultaneous small-ball sign assignments, but any improvement must exploit more structure than this second-moment calculation records.

Status: The metric, Haar-phase, second-moment, and random-sign statements above are self-contained paper proofs from report 6. The phase-zero count is a repackaging of `ExactCounting`. The finite positive-kernel bridge, exact weighted Fourier score, and integer-phase hit bound are [\[kernel-verified Lean\]](#) in `FejerKernelCertificate`; Haar integration and Borel–Cantelli are not claimed there. Neither almost-everywhere extinction nor fixed-mode equidistribution is specialized to ϑ .

68 Seventh and eighth reports: compact laws and approximation rigidity

Reports 7 and 8 were produced independently from the same inherited unified report. They are therefore not successive revisions of one another. This section audits each continuation on its own terms and then reconciles the nonoverlapping conclusions. Their directions are complementary. Report 7 asks what fixed measure-theoretic, profinite, and adelic observables can see under a hypothetical counterexample. Report 8 asks whether polynomial, rational, or Hermite–Padé approximation can exploit the exact integer lattice. Neither closes the conjecture: the first proves that every fixed compact factor is already maximally distributed, while the second proves that the most natural approximation scales lose their gain when denominators are cleared.

68.1 Difference groups, triangle–Haar laws, and profinite Beatty factors

Assume in this subsection that the conjecture fails and let $\beta = B + \alpha$ be its least nonintegral solution, where $B = \lfloor \beta \rfloor$ and $0 < \alpha < 1$. Write

$$M = 2^\beta = 2^a W, \quad A = 3^\beta = 3^b Z, \quad W > 1 \text{ odd}, \quad Z > 1, \quad 3 \nmid Z.$$

The inherited structure theorem gives $\mathcal{S} = \mathbb{N}_0 + \mathbb{N}_0\beta$ and unique coordinates. Taking differences produces an exact measure-theoretic dichotomy:

$$\mathcal{S} - \mathcal{S} = \begin{cases} \mathbb{Z}, & \text{if the conjecture holds,} \\ \mathbb{Z} + \mathbb{Z}\beta, & \text{if it fails.} \end{cases} \quad (265)$$

Modulo $\mathbb{Z}$, the failure branch is the countable dense Haar-null subgroup $\{k\alpha : k \in \mathbb{Z}\} \subset \mathbb{T}$. Thus closure, Haar measure, and the arithmetic set itself give three sharply different pictures of the same orbit. This is the first reason that a fixed weak-limit argument cannot settle the problem.

The macroscopic and microscopic coordinates can nevertheless be combined exactly. Put

$$L_T = \{(n, k) \in \mathbb{N}_0^2 : n + \beta k \leq T\}, \quad N_T = \#L_T = \frac{T^2}{2\beta} + O_\beta(T),$$

and $\Delta = \{(u, v) \in \mathbb{R}_{\geq 0}^2 : u + v \leq 1\}$. Report 7 proves that for every continuous $F : \Delta \times \mathbb{T} \rightarrow \mathbb{C}$,

$$\frac{1}{N_T} \sum_{(n,k) \in L_T} F\left(\frac{n}{T}, \frac{\beta k}{T}, \{k\beta\}\right) \longrightarrow 2 \int_\Delta \int_{\mathbb{T}} F(u, v, t) dt du dv. \quad (266)$$

For a nonzero circle Fourier mode, Abel summation reduces the assertion to bounded geometric partial sums of $e(hk\beta)$; the zero mode is an anisotropic Riemann sum. In the variables $H = u + v$ and $C = v/(u + v)$, the limiting density is $2H dH dC$, independently of the Haar phase. This is the precise independence statement; the separate Dirichlet description of $(u, v, 1 - u - v)$ is merely another description of uniform triangle area, not a replacement for the (H, C) factorization.

The finite congruence refinement is stronger. For every irrational α , every $d \geq 1$, residues $r, j \pmod{d}$, and interval $I \subset \mathbb{T}$,

$$\frac{1}{K} \#\{0 \leq k < K : k \equiv r \pmod{d}, \lfloor k\alpha \rfloor \equiv j \pmod{d}, \{k\alpha\} \in I\} \longrightarrow \frac{|I|}{d^2}. \quad (267)$$

Indeed, after writing $k = d\ell + r$, the map $x \mapsto (\lfloor dx \rfloor, \{dx\})$ sends Lebesgue measure to uniform measure on $\mathbb{Z}/d\mathbb{Z}$ times Lebesgue measure. Equivalently,

$$k \mapsto (k, \lfloor k\alpha \rfloor, \{k\alpha\})$$

is Haar distributed in $\widehat{\mathbb{Z}}^2 \times \mathbb{T}$. Every fixed family of congruence conditions is therefore asymptotically independent of the real phase. With the inherited finite irrationality measure $\|h\beta\|_{\mathbb{Z}} \geq c_M h^{1-\mu_M}$, where we harmlessly enlarge μ_M so that $\mu_M > 1$, the same argument controls growing moduli only in the range $d^{\mu_M+1} = o(K)$. The true denominator moduli grow exponentially in k , so this effective range remains far too small.

68.2 Exact denominator drift and noncompact escape

The compact arc coordinates

$$R_k = \frac{M^k}{2^{\lfloor k\beta \rfloor}}, \quad S_k = \frac{A^k}{3^{\lfloor k\beta \rfloor}}$$

separate into primitive numerators and denominator exponents

$$q_k^{(2)} = \lfloor k\beta \rfloor - ka, \quad q_k^{(3)} = \lfloor k\beta \rfloor - kb.$$

Writing $t_k = \{k\beta\}$ gives the exact identities

$$q_k^{(2)} = k(\beta - a) - t_k, \quad q_k^{(3)} = k(\beta - b) - t_k, \quad (268)$$

$$q_k^{(2)} - q_k^{(3)} = k(b - a). \quad (269)$$

Both centered valuation defects are therefore the same bounded quantity $-t_k$, rather than independent noise. Since $\beta - a = \log_2 W > 0$ and $\beta - b = \log_3 Z > 0$,

$$|R_k|_2 = 2^{q_k^{(2)}} \longrightarrow \infty, \quad |S_k|_3 = 3^{q_k^{(3)}} \longrightarrow \infty. \quad (270)$$

Thus the empirical measures escape every compact subset of $\mathbb{Q}_2 \times \mathbb{Q}_3$. Removing the denominator powers leaves (W^k, Z^k) in the compact group $\mathcal{H} = \overline{\langle (W, Z) \rangle} \subset \mathbb{Z}_2^\times \times \mathbb{Z}_3^\times$; the pair (t_k, W^k, Z^k) is distributed according to Lebesgue measure times Haar measure on $\mathcal{H}$. In a moving frame the limiting measure is the pushforward of $dt \, dm_{\mathcal{H}}$ under

$$(t, u, v) \mapsto (t, -t, -t, u, v).$$

This retains exactly the phase information that one-point compactification would erase.

The integer definition, the cast identity, the common centered defect, the cross-base difference, the natural-exponent layer, and the resulting 2- and 3-adic escape are [\[kernel-verified Lean\]](#) in `AdelicDrift`. The module deliberately defines the raw exponent in $\mathbb{Z}$ before proving nonnegativity, so no hidden natural-subtraction assumption enters the statement.

68.3 Sturmian rigidity and fixed-observable blindness

Choose a representative $t \in [0, 1)$ of an initial circle point and set

$$\xi_j(t) = \mathbf{1}_{[1-\alpha, 1)}(\{t + j\alpha\}).$$

Each bit is the carry in addition by α , so the exact telescoping identity is

$$\sum_{j=0}^{N-1} \xi_j(t) = \lfloor t + N\alpha \rfloor - \lfloor t \rfloor. \quad (271)$$

For uniform t , this sum takes only $\lfloor N\alpha \rfloor$ and $\lceil N\alpha \rceil$, with mean $N\alpha$ and variance $\{N\alpha\}(1 - \{N\alpha\}) \leq 1/4$. If $\delta_\ell = \{\ell\alpha\}$, then

$$\text{Cov}(\xi_0, \xi_\ell) = \max(0, \alpha - \delta_\ell) + \max(0, \alpha + \delta_\ell - 1) - \alpha^2. \quad (272)$$

The correlations do not decay. The process is a zero-entropy, pure-point factor of an irrational rotation; Bernoulli, Poisson, and central-limit heuristics are therefore the wrong model for its carries.

More generally, let G be a profinite group and let $\pi : \widehat{\mathbb{Z}}^2 \rightarrow G$ be continuous. Combining (266) and (267) shows that for continuous observables the scaled triangle coordinate, $\pi(k, \lfloor k\alpha \rfloor)$, and the circle phase have the product limit. The same holds for indicators whose discontinuity set has limiting measure zero; this qualification is needed for Sturmian bits, which are not continuous functions. Consequently every fixed compact congruence or unit factor is blind to the counterexample. A successful observable must move with height, retain an unbounded valuation, or resolve an exact shrinking target.

Two further paper-level calibrations reinforce the point. First,

$$\frac{1}{\log X} \int_1^X \delta_{(\{u\}, \{u^\vartheta\})} \frac{du}{u} \xrightarrow{*} m_{\mathbb{T}^2};$$

nonzero Fourier modes follow from a first-derivative estimate. Second, for $E_m(\psi) = \{c \in [a, b] : \|m^c\|_{\mathbb{Z}} < \psi_m\}$ one has $\lambda(E_m(\psi)) \ll_{a,b} \psi_m$, and for $\psi_m = m^{-\tau}$,

$$\dim_{\text{H}} \limsup_m E_m(m^{-\tau}) \leq \min\left(1, \frac{b+1}{b+\tau}\right). \quad (273)$$

This can place ϑ in a zero-dimensional exceptional set without excluding that named number. The local mass estimate near exact hits in report 7 is only a scale diagnostic unless overlaps and uniform Taylor remainders are controlled, and is not promoted here to a theorem. The near-Sidon statement for conditional truncations is likewise retained only as a conditional consequence of the [classical/open background] Evertse–Schlickewei–Schmidt finite-rank unit-equation theorem.

68.4 Arithmetic resolution of normalized polynomial approximation

Report 8 starts from a candidate $m \in [2^T, 2^{T+1})$, puts $u = m/2^T \in [1, 2)$ and $n = m^\vartheta = 3^T u^\vartheta \in \mathbb{N}$, and keeps careful track of the denominators introduced by normalization. If $p \in \mathbb{Q}[t]$, $\deg p \leq d$, and $Hp \in \mathbb{Z}[t]$, then

$$H2^{Td}(n - 3^T p(u)) \in \mathbb{Z}. \quad (274)$$

Hence an error smaller than $(H3^T 2^{Td})^{-1}$ forces exact equality at every candidate in the block. For integral A, B of degrees at most d , the correct rational-approximation object is the numerator residual, not the quotient error:

$$2^{Td}(B(u)n - 3^T A(u)) = 3^T 2^{Td}(B(u)u^\vartheta - A(u)) \in \mathbb{Z}. \quad (275)$$

A small quotient can therefore be arithmetically useless when B is large or has costly coefficients.

The decisive analytic comparison is elementary. For every $d \geq 1$ and every real polynomial p of degree at most d ,

$$\|t^\vartheta - p(t)\|_{[1,2]} \geq c_d := \frac{|(\vartheta)_{d+1}| 2^{\vartheta-2d-2}}{(d+1)^{d+1}}. \quad (276)$$

Indeed, with $h = (d+1)^{-1}$, the $(d+1)$ -st forward difference of p vanishes, the triangle inequality costs 2^{d+1} , and the iterated finite-difference mean-value theorem evaluates the same difference of t^ϑ at an intermediate point. For $d = 0$, endpoint separation gives the lower bound 1.

Combining (274) and (276) yields the all-degree obstruction:

Theorem 68.1 (Normalization-tax obstruction). *If $T \geq 4$, $d \geq 0$, $H \geq 1$, $\deg p \leq d$, and $Hp \in \mathbb{Z}[t]$, then*

$$H3^T 2^{Td} \|t^\vartheta - p(t)\|_{[1,2]} \geq 1. \quad (277)$$

Thus no normalized rational polynomial, in any degree, crosses the arithmetic forcing threshold.

The proof reduces to $T = 4, H = 1$. Writing $L_d = 3^4 2^{4d} c_d$, the rational bounds $3/2 < \vartheta < 8/5$ show $L_1 > \dots > L_5 > 4$, $L_6 > 1$, and $L_{d+1}/L_d > 1$ for $d \geq 6$; in particular the exact calculation includes $L_5 > 13328/3125$. No finite search bound is used. The forward-difference estimate, polynomial degree floor, exact denominator clearing, and the all-degree tax theorem are [\[kernel-verified Lean\]](#) in `PolynomialApproximationTax`; the sharp Bernstein-ellipse rate and its degree constant remain classical paper inputs.

Working in the physical variable avoids the factor 2^{Td} , but then the same lower bound requires degree

$$d \geq \frac{\log 3}{\log(4e)} T - O(\log T) = 0.460384\dots T - O(\log T).$$

Classical Bernstein approximation improves the slope to $\log 3 / \log(3 + 2\sqrt{2}) = 0.623238\dots$. Either slope exceeds the conditional block population slope by a wide margin; this is a structural mismatch, not a numerical failure at the beginning of the search range.

68.5 Rational contact rigidity and the area-to-diameter collapse

The local approximation audit has a sharper exact conclusion. For a polynomial P and a rational point (a, b) , let $\text{mult}_{(a,b)} P$ be the least total degree in the Taylor expansion of $P(a + U, b + V)$.

Theorem 68.2 (Rational contact equals algebraic multiplicity). *Let τ be transcendental, $0 \neq P \in \mathbb{Q}[X, Y]$, and $(a, b) \in \mathbb{Q}_{>0}^2$. If $s = \text{mult}_{(a,b)} P$, then*

$$\text{ord}_{z=0} P(ae^z, be^{\tau z}) = s. \quad (278)$$

If $H_s(U, V) = \sum_{j=0}^s h_j U^{s-j} V^j$ is the first nonzero homogeneous Taylor layer, the coefficient of z^s is $H_s(a, b\tau) = \sum_j h_j a^{s-j} b^j \tau^j$. This is a nonzero rational polynomial evaluated at a transcendental number, so it cannot vanish. In particular, for bidegree at most (d, e) the contact order is at most $d + e$; equality is attained only by a rational scalar multiple of $(X - a)^d (Y - b)^e$. The square of $(d + 1)(e + 1)$ coefficient slots thus collapses to the diameter $d + e$ over $\mathbb{Q}$. Over $\mathbb{Q}(\tau)$ the ordinary distinct-frequency Vandermonde instead permits order $(d + 1)(e + 1) - 1$.

For not-both-zero $A, B \in \mathbb{Q}[t]$ of degree at most d this gives

$$\text{ord}_{t=1} (B(t)t^\tau - A(t)) \leq d + 1.$$

If $A(1) = B(1)$ and $B(1) \neq 0$, the intersection is already transverse because

$$B(1)\tau + B'(1) - A'(1) \neq 0. \quad (279)$$

The rational-polynomial transcendence lemma, nonvanishing of the homogeneous leading coefficient, the derivative identity, and the Padé transversality endpoint are [\[kernel-verified Lean\]](#) in `RationalContactRigidity`. The general Taylor-order packaging in Theorem 68.2 remains a self-contained paper proof.

Finally, for a finite support $S \subset \mathbb{N}_0^2$ and nonzero integer coefficients,

$$F(x) = \sum_{(i,j) \in S} c_{ij} x^{i+j\vartheta}$$

is integral at every candidate. If $\|F\|_I < 1$, it vanishes at all candidates in I , while the accepted exponential-polynomial zero bound gives at most $|S| - 1$ positive zeros. Hence

$$\#(\mathcal{C} \cap I) \leq |S| - 1. \quad (280)$$

This arithmetic Müntz certificate is a direct corollary of the existing `ExponentialPolynomialZeros` machinery, not a new approximation theorem. The genuinely open target is to find an integral coefficient vector with norm below one after all archimedean, denominator, and p -adic height costs are charged. The schematic adelic height functional proposed in report 8 records this target but is not itself a proved estimate.

Status: Reports 7 and 8 are independent continuations from the common unified base. The exact drift/escape package is [\[kernel-verified Lean\]](#) in `AdelicDrift`; the finite-difference and normalization-tax package is [\[kernel-verified Lean\]](#) in `PolynomialApproximationTax`; the exact carry, prefix, mean, variance, joint-probability, and covariance package is [\[kernel-verified Lean\]](#) in `SturmianProbability`; and the homogeneous contact obstruction plus rational Padé transversality are [\[kernel-verified Lean\]](#) in `RationalContactRigidity`. The triangle–Haar, profinite Beatty, logarithmic Haar, Hausdorff-dimension, full contact-order, and area-to-diameter statements above have self-contained paper proofs. The effective discrepancy uses the inherited Baker–Matveev interface, the near-Sidon law uses the external finite-rank unit-equation theorem, and the sharp Bernstein rate is classical. The reported 2^{29} computation is not promoted: this unified report retains its reproducible 2^{27} trust frontier.

69 Ninth report: analytic common divisors and natural boundaries

Report 9 is another independent continuation from the common unified base. Its approximation section overlaps report 8 and is not repeated. Its distinct contribution is a three-part analytic reformulation: cancel the automatic integer zeros of the two sine compositions, identify the resulting conditional zero lattice and its canonical divisor, and encode the synchronized denominator exponents by integer power series with a natural boundary. A final compact-height count sharpens the earlier qualitative compact-orbit equivalence.

69.1 Removing the automatic zeros and counting the conditional divisor

Put

$$F_a(z) = \sin(\pi a^z), \quad a \in \{2, 3\}.$$

The accepted complex-localization theorem already shows that every common zero of F_2 and F_3 is real and simple. The nonnegative integers are automatic common zeros. They can be removed without introducing new ones by filling the removable singularities in

$$G_a(z) = \Gamma(-z) \sin(\pi a^z). \quad (281)$$

Indeed, at $n \in \mathbb{N}_0$ one has

$$G_a(n) = \frac{\pi(\log a)a^n}{n!}(-1)^{n+1+a^n} \neq 0, \quad (282)$$

so G_2 and G_3 have a common zero exactly when the conjecture has a nonintegral solution.

Under failure, let β be the least nonintegral solution. The exact conditional monoid classification identifies the de-integralized common zero set as

$$\Lambda_\beta = \{n + k\beta : n \in \mathbb{N}_0, k \in \mathbb{N}\}.$$

Its representation is unique, and elementary lattice counting gives

$$N_\beta(R) := \#(\Lambda_\beta \cap [0, R]) = \frac{R^2}{2\beta} + O_\beta(R). \quad (283)$$

Consequently its exponent of convergence is exactly two. With $E_2(w) = (1 - w) \exp(w + w^2/2)$, the genus-two product

$$P_\beta(z) = \prod_{\lambda \in \Lambda_\beta} E_2(z/\lambda) \quad (284)$$

converges locally uniformly, has precisely those simple zeros, and divides both G_2 and G_3 in the ring of entire functions. Conversely, any common nonunit entire divisor has a zero and hence gives a nonintegral solution. Thus the conjecture is equivalently the absence of a common nonunit entire divisor.

This is not an ordinary finite-order zero-counting problem: along suitable horizontal lines, $\log M(r, G_a) \geq c_a a^r$ for infinitely many r . The conditional divisor has only quadratic zero density inside functions of infinite order. It is also too large for a bounded-characteristic function on the right half-plane, since partial summation gives

$$\sum_{\lambda \in \Lambda_\beta} \frac{\lambda}{1 + \lambda^2} = \infty. \quad (285)$$

This supplies a precise but still open analytic target: a growth-controlled common divisor or Bézout theorem for the pair G_2, G_3 . Canonical-product convergence and the bounded-characteristic/Blaschke implication are classical analytic interfaces, not new Lean claims.

69.2 The synchronized word has a natural boundary

Write $\beta = B + \alpha$ with $0 < \alpha < 1$, and define

$$e_k = \lfloor k\alpha \rfloor, \quad s_k = e_{k+1} - e_k \in \{0, 1\}.$$

For the formal power series

$$\mathcal{S}(X) = \sum_{k \geq 0} s_k X^k, \quad \mathcal{E}(X) = \sum_{k \geq 1} e_k X^k,$$

telescoping gives the exact identity

$$X\mathcal{S}(X) = (1 - X)\mathcal{E}(X). \quad (286)$$

Both analytic series have radius one. The binary Sturmian word (s_k) is not eventually periodic because α is irrational. A rational finite-alphabet power series has an eventually periodic coefficient sequence; hence $\mathcal{S}$ is not rational. The classical Pólya–Carlson dichotomy now makes the unit circle a natural boundary for $\mathcal{S}$, and (286) transfers that boundary to $\mathcal{E}$.

The exact reduced denominator exponents have the form

$$q_{r,k} = rk + e_k, \quad \mathcal{Q}_r(X) = \frac{rX}{(1-X)^2} + \mathcal{E}(X).$$

Thus every $\mathcal{Q}_r$ has the same natural boundary, and

$$\mathcal{Q}_r - \mathcal{Q}_s = \frac{(r-s)X}{(1-X)^2}. \quad (287)$$

The carry telescope, formal-series identity, denominator decomposition, twice-cleared defect, and cross-base cancellation are [\[kernel-verified Lean\]](#) in `SturmianGeneratingSeries`. The natural-boundary conclusion itself invokes the classical Pólya–Carlson theorem. It is a one-slope diagnostic, not a contradiction: every irrational rotation already has this boundary, so a successful continuation theorem must use both prime supports together.

69.3 A logarithmic compact-height criterion

Let

$$\mathcal{C}_{2,3} = \{(2^t, 3^t) : 0 \leq t < 1\} \cap (\mathbb{Z}[1/2] \times \mathbb{Z}[1/3]),$$

and for a reduced rational pair define $H(a/b, c/d) = \max(|a|, b, |c|, d)$. Write $N_{2,3}(H)$ for the number of points of $\mathcal{C}_{2,3}$ of height at most H . The accepted compact-orbit theorem gives the qualitative equivalence $\mathcal{C}_{2,3} = \{(1, 1)\}$ if and only if the conjecture holds. Report 9 adds the exact growth dichotomy

$$N_{2,3}(H) = \begin{cases} 1, & \text{if the conjecture holds,} \\ \asymp_{\beta} \log H, & \text{if it fails.} \end{cases} \quad (288)$$

In the failure branch, the distinct points $(2^{\{k\alpha\}}, 3^{\{k\alpha\}})$ exhaust the localized arc, and their reduced numerator and denominator heights grow between two fixed exponentials in k . Therefore each of

$$N_{2,3}(H) = O(1), \quad N_{2,3}(H) = o(\log H)$$

is equivalent to the conjecture. Compact normalization has removed the automatic integer direction from the inherited $\asymp (\log H)^2$ unbounded count and leaves precisely one conditional generator.

Status: Report 9 was audited independently from the same common base as the other reports. The formal generating-series algebra is [\[kernel-verified Lean\]](#) in `SturmianGeneratingSeries`. The Gamma cancellation, quadratic lattice count, compact height dichotomy, and Blaschke calculation have self-contained paper proofs; canonical-product and Pólya–Carlson steps are classical interfaces. The duplicated Bernstein approximation material is represented once in Section 68. The inherited reported 2^{29} scan is not promoted; the unified trust frontier remains the reproducible 2^{27} computation.

70 Tenth report: matrix powers, Loewner certificates, and the one-half barrier

Report 10 is an independent matrix-theoretic continuation from the common base. It produces one exact reformulation and a family of genuine determinant certificates, but it also proves that polynomial representations of the original matrix carry no information beyond the two assumed integer outputs. The optimized determinant family reaches a limiting one-half height exponent and therefore remains on the same side of the four-exponentials wall.

70.1 One fixed integral matrix

Consider

$$A_0 = \begin{pmatrix} 0 & -6 \\ 1 & 5 \end{pmatrix}. \quad (289)$$

Its eigenvalues are 2, 3, and the complementary integral spectral projectors are $P_2 = 3I - A_0$ and $P_3 = A_0 - 2I$. Real functional calculus therefore gives

$$A_0^t = 2^t P_2 + 3^t P_3 = (3 \cdot 2^t - 2 \cdot 3^t)I + (3^t - 2^t)A_0. \quad (290)$$

Consequently

$$A_0^t \in M_2(\mathbb{Z}) \iff 2^t, 3^t \in \mathbb{Z}. \quad (291)$$

For the nontrivial direction, if $B = A_0^t$ is integral then $c = B_{21} = 3^t - 2^t$ and $d = B_{11} = 3 \cdot 2^t - 2 \cdot 3^t$ are integers, while $2^t = d + 2c$ and $3^t = d + 3c$. Equivalently,

$$\text{Alaoglu–Erdős} \iff \{t \in \mathbb{R} : A_0^t \in M_2(\mathbb{Z})\} = \mathbb{N}_0.$$

The concrete matrix, all four entry formulas, the exact entry-integrality equivalence, and this conjectural equivalence are [\[kernel-verified Lean\]](#) in `IntegralPowerMatrix`.

This reformulation is exact but intentionally not advertised as extra arithmetic. The unimodular matrix $\begin{pmatrix} -3 & -2 \\ 1 & 1 \end{pmatrix}$ diagonalizes A_0 over $\mathbb{Z}$, so the lattice already splits into its 2- and 3-eigensummands. Every polynomial, tensor, symmetric, or exterior-power construction from (A_0, A_0^t) is a polynomial in the two integers $(2^t, 3^t)$ and their monomials. It cannot recover the common real time unless it introduces analytic dependence, a denominator-sensitive inequality, or a new logarithmic direction.

70.2 Cross-Loewner and generalized-alternant certificates

For distinct positive nodes $z_1, \dots, z_N$ and distinct real exponents $\gamma_1, \dots, \gamma_N$, strict total positivity of e^{uv} gives

$$\det(z_i^{\gamma_j}) \neq 0,$$

with positive sign when both lists are increasing. If t is a nonintegral solution, the two-block exponent list

$$0, 1, \dots, r-1, \quad t, t+1, \dots, t+r-1 \quad (292)$$

is distinct. At $2r$ distinct nodes in the 2,3-smooth semigroup every entry is integral, so the alternant is a nonzero integer.

The same certificate admits an exact divided-difference form. For disjoint r -tuples $X = (x_i)$ and $Y = (y_j)$ and a function f on their union, put

$$L_f(X, Y)_{ij} = \frac{f(x_i) - f(y_j)}{x_i - y_j}.$$

If $M_f(X, Y)$ is the $2r$ -row alternant with columns $1, z, \dots, z^{r-1}, f(z), zf(z), \dots, z^{r-1}f(z)$, then

$$\det M_f(X, Y) = (-1)^{r(r+1)/2} \left(\prod_{i,j} (x_i - y_j) \right) \det L_f(X, Y). \quad (293)$$

For $f(u) = u^t$ and smooth integer nodes, the right-hand side is therefore a nonzero integer. Writing $t = k + \alpha$ with $0 < \alpha < 1$, the fractional-power kernel

$$K_\alpha(u, v) = \frac{u^\alpha - v^\alpha}{u - v}$$

has the positive Stieltjes representation

$$K_\alpha(u, v) = \frac{\sin(\pi\alpha)}{\pi} \int_0^\infty \frac{\lambda^\alpha}{(u + \lambda)(v + \lambda)} d\lambda.$$

Thus its ordered cross minors are strictly positive. Combining positivity with (293) gives the explicit positive integer

$$\left(\prod_i x_i^k y_i^k \right) \left| \prod_{i,j} (x_i - y_j) \right| \det(K_\alpha(x_i, y_j)) \in \mathbb{Z}_{>0}. \quad (294)$$

The positivity holds for every $0 < \alpha < 1$; all arithmetic content is in the displayed normalization. These finite determinant identities have paper proofs here. They are not silently conflated with the already formal generalized-Vandermonde positivity theorem.

70.3 Quantitative repulsion and the optimized barrier

Let $z_i = e^{u_i}$ with pairwise distinct $u_i \in [T, T + \delta]$, $0 < \delta \leq 1$, and pairwise distinct exponents $0 \leq \gamma_j \leq B$. (With repeated nodes or exponents the determinant vanishes and the upper bound is trivial.) Put $\Sigma = \sum_j \gamma_j$, $D = N(N-1)/2$, and $B_* = \max(1, B)$. Newton divided differences followed by Hadamard's inequality give the explicit upper bound

$$\left| \det(z_i^{\gamma_j}) \right| \leq e^{T\Sigma} \delta^D \frac{N^{N/2} e^{NB} B_*^D}{\prod_{\ell=0}^{N-1} \ell!}. \quad (295)$$

Together with the nonzero-integer lower bound, this excludes a sufficiently tight cluster of $2r$ smooth nodes. For the equal two-block list (292) it gives an explicit bound of the shape

$$\delta \geq C_{r,t}^{-1/[r(2r-1)]} R^{-(t+r-1)/(2r-1)}.$$

The exponent tends to $1/2$ as $r \rightarrow \infty$, much smaller than the logarithmic spacings available from counting smooth numbers.

Allowing unequal block sizes does not improve the limiting ratio. With $p + q = N$ and exponent set $0, \dots, p-1, t, \dots, t+q-1$, the total exponent is

$$\Sigma(p, q; t) = \frac{p(p-1)}{2} + qt + \frac{q(q-1)}{2} = D + q(q+t-N), \quad (296)$$

so

$$\min_{p+q=N} \frac{\Sigma(p, q; t)}{D} = \frac{1}{2} + O_t(N^{-1}). \quad (297)$$

Initial consecutive shifts already minimize the height inside each coset. Translating signed exponents into the first quadrant restores exactly the denominator height that cancellation appeared to save. This is the matrix one-half barrier.

There is a dual exact identity for candidate exponents. Let $d \geq 1$, let $p, q \in \mathbb{N}$ with $q > 0$, and put $\varepsilon = qt - p$ and $\tau_i = (d-1-i)p + iqt$ for $0 \leq i < d$. Then at distinct smooth bases s_j , $0 \leq j < d$,

$$\det(s_j^{\tau_i}) = \left(\prod_j s_j^{(d-1)p} \right) \prod_{j < k} (s_k^\varepsilon - s_j^\varepsilon). \quad (298)$$

Its nonzero-integral lower bound yields an explicit but exponentially weak irrationality estimate for $|qt - p|$. Higher dimension organizes the elementary integer $2^{qt} - 2^p$ but does not change its exponential scale. The affine-exponent determinant, its real-power specialization, the exact semigroup exponent identity, and (298) are [\[kernel-verified Lean\]](#) in `CandidateAlternant`; the mean-value upper bound and resulting irrationality estimate remain the paper layer.

Status: Report 10 was audited independently from the common base. The exact fixed matrix reformulation is [\[kernel-verified Lean\]](#) in `IntegralPowerMatrix`. Generalized-Vandermonde nonvanishing and ordered-chamber positivity are separately [\[kernel-verified Lean\]](#) in the inherited determinant modules, and the exact rational-approximation alternant is [\[kernel-verified Lean\]](#) in `CandidateAlternant`. The Sylvester–Loewner factorization, positive normalized determinant, explicit alternant bound, and optimized one-half barrier are retained as paper theorems. No claim is made that formal representation theory manufactures a third exponential, or that matrix positivity by itself selects integral exponents.

71 Eleventh report: finite-place transversality and local rank-drop certificates

Report 11 is an independent p -adic and matrix-coefficient continuation from the common base. Its real prime-symbol polynomial, rational matrix-coefficient obstruction, and sharp support bound

overlap the already reconciled reports 1–4. They are represented once by **FormalPrimeDeterminant**, **RestrictedMatrixCoefficient**, and **SparsePowerCurve**. The genuinely distinct layer compares the singular real logarithm matrix with its finite-place analogues, isolates the first Fermat-quotient obstruction, and shows why ordinary real rational approximation does not propagate to a local rank drop.

71.1 The structural primes are locally transverse

For the Iwasawa logarithm $\log_p : \mathbb{Q}_p^* \rightarrow \mathbb{Q}_p$, normalized by $\log_p p = 0$, define

$$L_p(M, A) = \begin{pmatrix} \log_p 2 & \log_p M \\ \log_p 3 & \log_p A \end{pmatrix}, \quad D_p(M, A) = \log_p M \log_p 3 - \log_p A \log_p 2. \quad (299)$$

On positive rationals the kernel is exactly

$$\log_p u = 0 \iff u = p^k \text{ for some } k \in \mathbb{Z}. \quad (300)$$

Hence a nonintegral counterexample pair $(M, A) = (2^\beta, 3^\beta)$ satisfies

$$D_2(M, A) = \log_2 M \log_2 3 \neq 0, \quad (301)$$

$$D_3(M, A) = -\log_3 A \log_3 2 \neq 0. \quad (302)$$

The real matrix has rank one, but both finite-place matrices at the primes defining the problem have rank two. If, for a prime p , the pair (M, A) belonged to the closure of $\{(2^n, 3^n) : n \in \mathbb{Z}\}$ in $(\mathbb{Q}_p^*)^2$, continuity of $\log_p$ would instead force $D_p(M, A) = 0$. Thus a real-to-local closure principle at either structural prime would prove the conjecture, but no such transfer theorem is known.

This is not merely a change of language. It reverses the naive local expectation: a hypothetical counterexample is singular only at the real place among the three places that are directly visible from its definition. Any adelic proof must compare a real zero with proven local nonzeros, not copy the real relation into $\mathbb{Q}_2$ or $\mathbb{Q}_3$.

71.2 The Fermat-quotient first layer

For an odd prime $p \nmid u$, set

$$q_p(u) = \frac{u^{p-1} - 1}{p} \pmod{p}.$$

The first Iwasawa-logarithm digit is

$$\frac{\log_p u}{p} \equiv -q_p(u) \pmod{p}. \quad (303)$$

Therefore, for $p \geq 5$ and $p \nmid 6MA$, the determinant

$$\Delta_p(M, A) = q_p(M)q_p(3) - q_p(A)q_p(2) \pmod{p} \quad (304)$$

satisfies

$$\frac{D_p(M, A)}{p^2} \equiv \Delta_p(M, A) \pmod{p}. \quad (305)$$

Thus $\Delta_p \neq 0$ certifies the exact first possible valuation $v_p(D_p) = 2$. Conversely, for every integer-exponent pair $(M, A) = (2^n, 3^n)$,

$$q_p(u^n) \equiv nq_p(u) \pmod{p}$$

makes $\Delta_p(M, A) = 0$. A local rank drop must therefore pass a generalized Wieferich congruence and then an infinite tower of higher congruences. The finite near-miss scan in report 11 is retained only as a software diagnostic; it is not evidence that every primitive pair has a bounded detecting prime.

71.3 Real approximants and the exact adelic wedge

Let r/s be a rational approximation to β , $s > 0$, and put

$$R_2(r, s) = \frac{M^s}{2^r}, \quad R_3(r, s) = \frac{A^s}{3^r}.$$

At the real place the errors are perfectly synchronized:

$$\log R_2 = (s\beta - r) \log 2, \quad \log R_3 = (s\beta - r) \log 3.$$

At a finite place set

$$E_p(r, s) = (s \log_p M - r \log_p 2, s \log_p A - r \log_p 3).$$

Direct expansion gives the exact wedge identity

$$E_{p,1}(r, s) \log_p 3 - E_{p,2}(r, s) \log_p 2 = sD_p(M, A). \quad (306)$$

At the structural primes the numerator r disappears from one coordinate:

$$\log_2 R_2(r, s) = s \log_2 M, \quad v_2(\log_2 R_2) = v_2(s) + v_2(\log_2 M), \quad (307)$$

$$\log_3 R_3(r, s) = s \log_3 A, \quad v_3(\log_3 R_3) = v_3(s) + v_3(\log_3 A). \quad (308)$$

Choosing a better continued-fraction numerator improves the real approximation but has no effect on these two local depths; only divisibility of the denominator s helps. This exact anti-transfer explains why an archimedean convergent does not automatically produce a local near-singularity.

71.4 Cascade depth and the Kummer norm diagnostic

The report also records two useful no-go estimates. First, suppose each of the first $J(T)$ active tropical layers contributes at most $CT^{3/2}$ fixed-base logarithmic factors, each of valuation at most $C'(\log T)^2$. Their total exposed valuation is then

$$O\left(J(T)T^{3/2}(\log T)^2\right). \quad (309)$$

Against $H_T \asymp T^{5/2}$ this is lower order whenever $J(T) = o(T/(\log T)^2)$. This is a conditional bookkeeping proposition for a factorized layer model, not an upper bound on arbitrary determinant cancellation; a global higher initial form may behave differently.

Second, for positive real N th roots define the algebraic integer

$$\Xi_N = (M^{1/N} - 1)(3^{1/N} - 1) - (A^{1/N} - 1)(2^{1/N} - 1).$$

The determinant relation cancels the quadratic term and leaves

$$\Xi_N = \frac{\beta \log 2 \log 3 (\beta - 1)(\log 2 - \log 3)}{2N^3} + O_{M,A}(N^{-4}). \quad (310)$$

Thus the distinguished embedding is polynomially small and nonzero for large N . The field degree and varying, generally uncontrolled root-of-unity twists (subject to the algebraic relations among the radicands) grow with N , however, so a crude product-formula bound becomes worse rather than better. A viable algebraic-integer construction would need simultaneous smallness in many conjugates or exponential smallness at one place.

Status: Report 11 was audited independently from the common base. Its real prime-symbol and matrix-coefficient layers are already [\[kernel-verified Lean\]](#) in `FormalPrimeDeterminant` and `RestrictedMatrixCoefficient`; its sparse-grid core is [\[kernel-verified Lean\]](#) in `SparsePowerCurve`. The exact Fermat-quotient divisibility identity, power law modulo p , determinant interface, and its vanishing on $(2^n, 3^n)$ are [\[kernel-verified Lean\]](#) in `FermatQuotientDeterminant`. The passage from that residue determinant to the Iwasawa logarithm uses the classical congruence (303); the structural-prime theorem uses the classical Iwasawa-logarithm kernel. The exact wedge, conditional cascade-depth, and Kummer expansion remain elementary paper results. No p -adic Matrix Coefficient or real-to-local transfer theorem is asserted.

72 Twelfth report: prime switches and microscopic resonance

Report 12 is another independent continuation from the common unified base. It returns to the historical origin of the problem: simultaneous changes in the prime exponents of colossally abundant numbers. This is a valuable structured test family, but its logical scope is narrower than the integer-output conjecture. A switch produces rational values p^ε and q^ε ; excluding all switches would prove only a special case of the stronger rational-output statement. No reduction from arbitrary rational outputs, or from this switch family, to Conjecture 2.1 is known.

72.1 Exact switch heights and logarithmic means

For a prime p and an integer $a \geq 1$, set

$$H_{p,a} = p + p^2 + \cdots + p^a, \quad \varepsilon_{p,a} = \frac{\log(1 + H_{p,a}^{-1})}{\log p}. \quad (311)$$

Comparing the two consecutive local factors in $\sigma(n)/n^{1+\varepsilon}$ gives the exact tie identity

$$p^{\varepsilon_{p,a}} = \frac{p^{a+1} - 1}{p^{a+1} - p} = 1 + \frac{1}{H_{p,a}}. \quad (312)$$

For fixed a this parameter decreases with p , and for fixed p it decreases with a . If $p < q$ and $\varepsilon_{p,a} = \varepsilon_{q,b}$, then

$$H_{p,a} > H_{q,b}, \quad a > b. \quad (313)$$

Thus simultaneous switches form an antichain in prime size and local exponent.

The analytic center of the calculation is the consecutive logarithmic mean

$$R(H) = \frac{1}{\log(1 + 1/H)}, \quad J(H) = H + \frac{1}{2} - \frac{1}{12H}.$$

For every real $H \geq 1$ one has the globally valid two-sided enclosure

$$J(H) < R(H) < J(H) + \frac{1}{24H^2}. \quad (314)$$

The proof differentiates two rational majorants for $\log(1 + x)$ on $0 < x \leq 1$; it is an inequality with an exact remainder, not merely an asymptotic expansion. The definitions and the complete enclosure (314) are [\[kernel-verified Lean\]](#) in `ConsecutiveLogMeanBounds`.

At a collision, $R(H_{p,a}) \log p = R(H_{q,b}) \log q$. Writing $H = H_{p,a}$ and $K = H_{q,b}$, (314) therefore gives

$$|J(H) \log p - J(K) \log q| < \frac{\log p}{24H^2} + \frac{\log q}{24K^2}. \quad (315)$$

This is a highly structured Dirichlet-scale approximation: the coefficients are geometric sums, and exploiting that structure rather than merely invoking a general irrationality measure is essential.

72.2 A microscopic prescribed-prime window

For fixed $b \geq 1$, put

$$H_b(t) = t + t^2 + \cdots + t^b, \quad \Phi_b(t) = J(H_b(t)) \log t.$$

Since $\Phi'_b(t) > bt^{b-1} \log t$ on $t > 1$, there is a unique real $Q_{p,a,b} > 1$ satisfying

$$\Phi_b(Q_{p,a,b}) = J(H_{p,a}) \log p.$$

If primes $p < q$ switch simultaneously, the mean-value theorem and (315) imply

$$|q - Q_{p,a,b}| < \frac{\frac{\log p}{24H_{p,a}^2} + \frac{\log q}{24H_{q,b}^2}}{b(q - \frac{1}{2})^{b-1} \log(q - \frac{1}{2})} = O_b(q^{1-3b}). \quad (316)$$

The interval has length less than one. Already for $b = 1$ its radius is $O(q^{-2})$. Consequently modern prime-in-short-interval theorems, whose intervals still have positive power length, cannot decide the switch: before primality enters, the nearest integer has already been prescribed. The report's suggested modernization of the 1944 prime-block argument is retained only as a schematic transference of known prime-distribution input; its required uniform error estimates are not reconstructed here. Likewise, the first-order formula for the location of $Q_{p,a,b}$ is a useful calibration, but its displayed proof in the source report omits a bootstrap and is not used as a theorem in this synthesis.

72.3 The exact 2–3 convergent sieve

For a possible switch between 2 at level a and 3 at level b , define

$$P = 2^{a+2} - 3, \quad Q = 3^{b+1} - 2. \quad (317)$$

The second-order enclosure, together with the elementary bounds $3/2 < \vartheta < 8/5$, yields the exact necessary condition

$$0 < \vartheta - \frac{P}{Q} < \frac{1}{2Q^2}. \quad (318)$$

After reduction, P/Q must therefore be a continued-fraction convergent of ϑ . Equivalently, a collision must hit an exponentially shrinking shifted-power target

$$0 < \vartheta(3^{b+1} - 2) - (2^{a+2} - 3) < \frac{1}{2(3^{b+1} - 2)}.$$

This converts the transcendental equality into a discrete S -unit sieve, but no uniform argument over all convergents is currently known.

The literal exact-integer verifier in report 12 checks this necessary condition for $1 \leq b \leq 10000$. It certifies 10000 unambiguous floors, finds the sole power-shape hit $(a, b, P, Q) = (5, 3, 125, 79)$, and rejects that hit by the stronger Legendre window; hence there are no survivors in the tested range. The verifier source has SHA-256

a7aef69a62ce5b07bd8f909966244b3a5e92e23ac9bb2e1b951433d1f6902ce8.

An independent rerun reproduced every mathematical field. The source report’s JSON digest is not used as a certificate because its payload includes nondeterministic elapsed time. This remains [\[software computation\]](#), not a kernel proof and not a finite test of the original integer-output conjecture.

72.4 What prime distribution can and cannot force

For fixed integers $U, V \geq 2$, define the universal Fermat-synchronized divisor

$$F_{U,V}(n) = \prod_{\substack{\ell \text{ prime, } \ell \nmid UV \\ \ell-1 \mid n}} \ell.$$

Fermat’s theorem gives

$$F_{U,V}(n) \mid \gcd(U^n - 1, V^n - 1), \quad \log F_{U,V}(n) \leq \tau(n) \log(n+1) = n^{o(1)}. \quad (319)$$

Along $n_y = \text{lcm}(1, \dots, \lfloor y \rfloor)$, the prime number theorem and the standard divisor estimate still leave $\log F_{U,V}(n_y) = n_y^{o(1)} = o(n_y)$. This is far below the exponential common-divisor growth that would contradict the Bugeaud–Corvaja–Zannier bound for multiplicatively independent U, V . Generic prime distribution therefore cannot manufacture the missing real-to-finite-place correlation. It becomes relevant only after another argument has produced a rigid prescribed integer, such as the nearest integer in (316) or a convergent satisfying (318).

Status: Report 12 was audited independently from the same common base as all other numbered reports. The exact logarithmic-mean enclosure is [\[kernel-verified Lean\]](#) in `ConsecutiveLogMeanBounds`. The switch formula, antichain, resonance, microscopic window, convergent sieve, and Fermat-synchronization barrier are self-contained paper arguments. The low-level switch exclusion is [\[software computation\]](#) with the scope and software qualification stated above. Modern short-interval results are scale calibrations only, and the claimed transference to the complete 1944 block argument is not promoted to a theorem. Nothing in this section reduces the rational switch problem to the original integer-output conjecture.

73 Thirteenth report: exterior powers and boundary realizations

Report 13 is independent of report 10 and of every other numbered continuation: it starts from the same inherited unified base and pursues a different matrix-analysis program. Its new content is not another reformulation of the Matrix Coefficient principle. It asks what finite-dimensional matrix phenomena a hypothetical counterexample can realize, whether proper exterior powers improve the determinant budget, and whether a K^2 -point multiplicative grid can be compressed to its $O(K)$ boundary data.

Throughout this section assume hypothetically that

$$M = 2^\beta \in \mathbb{Z}_{>0}, \quad A = 3^\beta \in \mathbb{Z}_{>0}, \quad \beta \notin \mathbb{Q}, \quad (320)$$

and retain $\vartheta = \log_2 3$. Then

$$G = \{2^r M^s : r, s \in \mathbb{Z}\} \subset \mathbb{Q}_{>0} \quad (321)$$

is dense in $\mathbb{R}_{>0}$ and every $g \in G$ has rational ϑ -th power.

73.1 Rational conditioning does not create rational rank

The singular logarithm matrix factors as

$$L = \begin{pmatrix} \log 2 & \log M \\ \log 3 & \log A \end{pmatrix} = \begin{pmatrix} \log 2 \\ \log 3 \end{pmatrix} (1, \beta).$$

For a nonzero $c \in \mathbb{R}^2$ and $B \geq 1$, set

$$\mu_c(B) = \min_{0 \neq z \in \mathbb{Z}^2, \|z\|_\infty \leq B} |z^\top c|, \quad \eta_L(B) = \min_{\substack{0 \neq u, v \in \mathbb{Z}^2 \\ \|u\|_\infty, \|v\|_\infty \leq B}} |u^\top L v|.$$

The rank-one factorization gives the exact identity

$$\eta_L(B) = \mu_{(\log 2, \log 3)}(B) \mu_{(1, \beta)}(B). \quad (322)$$

Dirichlet approximation supplies $\eta_L(B) \ll_M B^{-2}$, while effective linear forms in logarithms give only a computable polynomial lower bound. Approximate zero coefficients after bounded rational basis changes are therefore plentiful and entirely compatible with exact nonvanishing. The missing conclusion remains rational orthogonality in one of the two factors, exactly as in `RestrictedMatrixCoefficient`.

73.2 Projective arithmetic density and a Hadamard-power no-go

The density in (321) has a simultaneous finite-matrix strengthening. Given $n \geq 1$, a positive-entry real matrix B , and $\epsilon > 0$, choose each entry from G within a sufficiently small tolerance and clear all negative exponents by one common $\lambda = 2^R M^S$ with $R, S \geq 0$. Then $\lambda, \lambda^\vartheta \in \mathbb{Z}_{>0}$, and this produces

$$C \in M_n(\mathbb{Z}_{>0}), \quad C^{\circ\vartheta} \in M_n(\mathbb{Z}_{>0}), \quad (323)$$

such that

$$\|\lambda^{-1}C - B\|_{\max} < \epsilon, \quad \|\lambda^{-\vartheta}C^{\circ\vartheta} - B^{\circ\vartheta}\|_{\max} < \epsilon. \quad (324)$$

The choices may be made symmetrically. More precisely, if $\mathcal{U}$ and $\mathcal{V}$ are separately open and invariant under positive scalar multiplication, then $B \in \mathcal{U}$ and $B^{\circ\vartheta} \in \mathcal{V}$ imply that the integral pair may be chosen with $C \in \mathcal{U}$ and $C^{\circ\vartheta} \in \mathcal{V}$.

This retires a natural positivity route. For distinct positive $r_1, \dots, r_4$ and small $t > 0$, Cauchy–Binet gives

$$\det((1 + tr_i r_j)^\alpha)_{i,j=1}^4 = t^6 \left(\prod_{k=0}^3 \binom{\alpha}{k} \right) \prod_{i < j} (r_j - r_i)^2 + O(t^7). \quad (325)$$

For $1 < \alpha < 2$ the leading coefficient is negative. After a small positive diagonal perturbation, the unpowered matrix is positive definite while its entrywise α -power is indefinite. Applying (324) with $\alpha = \vartheta$ yields a positive-definite integer matrix C for which $C^{\circ\vartheta}$ is both integral and indefinite. Thus critical-exponent failure is compatible with a counterexample; it cannot serve as an integrality contradiction.

Supports placed on a rational near-kernel direction give no hidden collective gain either. Their alternants factor into an ordinary Vandermonde times a Schur polynomial; for consecutive support the determinant is exactly a product of pairwise integer binomials. This is the homogeneous counterpart of the affine factorization already [\[kernel-verified Lean\]](#) in `CandidateAlternant`, so it is cross-referenced rather than counted as a separate formal advance.

73.3 Determinantal divisors and the compound ceiling

Let $E \in M_N(\mathbb{Z})$ be nonsingular and $1 \leq k \leq N$. Write $d_k(E)$ for the gcd of its k -minors, let $\sigma_1(E) \geq \dots \geq \sigma_N(E) > 0$ be its singular values, adopt $v_p(0) = +\infty$, and define

$$\tau_{p,k}(E) = \min_{\substack{|I|=|J|=k \\ \sigma: I \xrightarrow{\sim} J}} \sum_{i \in I} v_p(E_{i,\sigma(i)}), \quad Q_k(E) = \prod_p p^{\tau_{p,k}(E)}. \quad (326)$$

Only finitely many primes contribute to $Q_k(E)$. The exact adelic compound sandwich is

$$Q_k(E) \mid d_k(E), \quad d_k(E) \leq \|\wedge^k E\|_2 = \prod_{j=1}^k \sigma_j(E). \quad (327)$$

The divisibility uses only termwise prime divisibility of every minor; the inequality views the minors as the entries of $\wedge^k E$. It is the entire Smith/exterior-power analogue of the inherited full-determinant comparison.

On the complete triangular evaluation grid E_T , put $Q_{T,k} = Q_k(E_T)$ and let $E_{T,k}^-$ and $E_{T,k}^+$ denote, respectively, the minimum reverse-pairing and maximum same-order k -point assignment energies,

and take $k/N_T \rightarrow \rho \in (0, 1]$. The limiting quantile is $\sqrt{t}$, so

$$\frac{E_{T,k}^-}{N_T T L_T \log 2} \rightarrow \int_0^\rho \sqrt{t(\rho - t)} dt = \frac{\pi}{8} \rho^2, \quad \frac{E_{T,k}^+}{N_T T L_T \log 2} \rightarrow \int_{1-\rho}^1 t dt = \rho - \frac{\rho^2}{2}. \quad (328)$$

The combinatorial factor $\binom{N_T}{k} k!$ is lower order on the T^4 scale. Under these specific termwise-divisor and largest-minor/Frobenius bounds one obtains the ceiling

$$\limsup_{T \rightarrow \infty} \frac{\log Q_{T,k}}{E_{T,k}^+ + \log \left(\binom{N_T}{k} k! \right)} \leq R(\rho) = \frac{\pi \rho}{8 - 4\rho}, \quad R(\rho) < R(1) = \frac{\pi}{4} \quad (0 < \rho < 1). \quad (329)$$

Thus every proper compound is worse than the full determinant under the inherited estimates. This is not a no-go theorem for exterior powers themselves: a sharper singular-value free energy, cancellation among minors, or extra Smith divisibility could still change the comparison.

The report also applies a two-variable finite-difference transform to the square-grid evaluation matrix. It extracts diagonal weights $a!b!$. If, for each prime p , the multiset $\{v_p(a!b!) : 0 \leq a, b < K\}$ is arranged as $\omega_{p,1} \leq \dots \leq \omega_{p,K^2}$ and

$$\mathfrak{F}_{K,k} = \prod_p p^{\sum_{\ell=1}^k \omega_{p,\ell}},$$

then the every-compound divisor is

$$\mathfrak{F}_{K,k} \mid d_k(\mathcal{E}_K) \quad (1 \leq k \leq K^2). \quad (330)$$

At top degree this becomes

$$\left(\prod_{r=0}^{K-1} r! \right)^{2K} \mid \det \mathcal{E}_K, \quad \log \left(\prod_{r=0}^{K-1} r! \right)^{2K} = K^3 \log K - \frac{3}{2} K^3 + O(K^2 \log K).$$

The mechanism is the two-variable/all-compound extension of `SuperfactorialDeterminant`, `AllLayerUnitBasis`, and `RowBlockUnitDivisibility`; its contribution is lower order than the K^4 determinant budget and is retained as a paper specialization, not advertised as a new leading divisor.

73.4 Exact boundary displacement and transfer determinants

Let $K > 0$, $\mathcal{I}_K = \{0, \dots, K-1\}^2$, and set

$$\mathcal{E}_K[(i, j), (u, v)] = (2^u M^v)^i (3^u A^v)^j. \quad (331)$$

Put $d_{2;u,v} = 2^u M^v$ and $d_{3;u,v} = 3^u A^v$. Let D_2, D_3 be the diagonal matrices with these respective entries, let S_2, S_3 be the two truncated forward row shifts, and let P_2, P_3 project onto their terminal row boundaries. Directly shifting a row index gives

$$S_2 \mathcal{E}_K - \mathcal{E}_K D_2 = -P_2 \mathcal{E}_K D_2, \quad S_3 \mathcal{E}_K - \mathcal{E}_K D_3 = -P_3 \mathcal{E}_K D_3. \quad (332)$$

After sorting, irrationality of ϑ and β , respectively, makes the row parameters and column frequencies pairwise distinct; strict total positivity of e^{xy} therefore gives $\det \mathcal{E}_K \neq 0$. Since the coordinate

multipliers are nonzero, each right side in (332) has rank exactly K . The commuting multiplication operators

$$T_r = \mathcal{E}_K D_r \mathcal{E}_K^{-1} \quad (r = 2, 3)$$

have the boundary-feedback form $T_r = S_r + R_r$ with rank $R_r = K$. Their commutativity gives the exact boundary relation

$$[S_2, R_3] + [R_2, S_3] + [R_2, R_3] = 0. \quad (333)$$

If $J_r : \mathbb{R}^K \rightarrow \mathbb{R}^{K^2}$ injects the corresponding boundary and $R_r = J_r C_r$, the matrix determinant lemma compresses the full spectrum, for $z \neq 0$, to

$$\prod_{0 \leq u, v < K} (z - d_{r;u,v}) = z^{K^2} \det \left(I_K - C_r (zI - S_r)^{-1} J_r \right), \quad (zI - S_r)^{-1} = \sum_{h=0}^{K-1} z^{-h-1} S_r^h. \quad (334)$$

This is an exact $K \times K$ encoding of a K^2 -point spectrum. It does not by itself provide an arithmetic gain: the primitive denominators and local Smith factors of the boundary feedback remain unknown.

Because the $d_{2;u,v}$ are distinct, there is also a unique $\widehat{Q}_K \in \mathbb{Q}[X]$ of degree below K^2 with $\widehat{Q}_K(d_{2;u,v}) = d_{3;u,v}$. Simultaneous diagonalization gives the exact paper identity

$$T_3 = \widehat{Q}_K(T_2). \quad (335)$$

Together with (333), this concentrates the unresolved arithmetic in two rank- K boundary feedbacks.

73.5 Kraus–Loewner positivity and its denominator tax

For $f(t) = t^\vartheta$ and $a, x, y > 0$, let $K_a(x, y) = f[x, a, y]$ be the second divided difference. The operator-convex integral representation gives

$$K_a(x, y) = c_\vartheta \int_0^\infty \frac{s^\vartheta}{(x+s)(a+s)(y+s)} ds, \quad c_\vartheta = \frac{\sin(\pi(\vartheta-1))}{\pi} > 0. \quad (336)$$

Let $0 < x_1 < \dots < x_n$ and $0 < y_1 < \dots < y_n$. Andréief's identity and the two Cauchy determinants give a strictly positive multiple integral for $\det(K_a(x_i, y_j))$. Thus the cross-kernel is strictly totally positive.

When a, x_i, y_j are pairwise distinct positive integers and all their ϑ -th powers are integers, each entry is rational and

$$\left(\prod_i (x_i - a) \right) \left(\prod_j (a - y_j) \right) \left(\prod_{i,j} (x_i - y_j) \right) \det(K_a(x_i, y_j)) \in \mathbb{Z} \setminus \{0\} \quad (337)$$

But scaling all nodes by $R > 0$ multiplies this cleared expression by $R^{n^2+n\vartheta}$, a positive power. (It remains an integer after scaling whenever the scaled nodes satisfy the same integrality hypotheses.) The derivative decay is therefore overwhelmed by the naive denominator; a useful argument would need primitive denominator cancellation or new numerator divisibility.

Status: Report 13 was audited independently from the common base, not as a revision of report 10. Its projective density and Hadamard-power no-go, compound sandwich and triangular ceiling, every-compound finite-difference specialization, transfer determinant, rational-conditioning factorization, and Kraus–Loewner package have self-contained paper proofs. The exact shift/displacement identities are [\[kernel-verified Lean\]](#) in **BoundaryDisplacement**; its rank- K endpoints explicitly assume $K > 0$, nonsingularity, and nonzero coordinate multipliers, all supplied for the displayed grid by the paper argument above. The transfer determinant and spectral interpolation remain the paper layer. The rational-direction alternant and factorial-basis mechanisms overlap existing accepted modules and are not double-counted. The compound ceiling applies only to the stated termwise-divisor/Frobenius estimates; the boundary compression and Kraus positivity are exact interfaces, not claimed contradictions.

74 Fourteenth report: structural-prime residuals and dilation results

Report 14 is another independent continuation from the common unified base. Its first finite-place observations overlap report 11: the real logarithm matrix has rank one, whereas the Iwasawa logarithm matrices at the structural primes 2 and 3 have rank two; the first good-prime shadow is the Fermat-quotient determinant already isolated in **FermatQuotientDeterminant**. The genuinely new part begins when the square-grid interpolant is placed in its integral coefficient lattice and its two boundary residuals are normalized at a structural prime.

74.1 Local rank is thick, but it is not a product-formula input

Under a hypothetical nonintegral solution write

$$M = 2^\beta \in \mathbb{N}, \quad A = 3^\beta \in \mathbb{N}, \quad \beta \notin \mathbb{Q}.$$

For Iwasawa logarithms the positive-rational kernel is

$$\log_p r = 0 \iff r \in p^{\mathbb{Z}} \quad (r \in \mathbb{Q}_{>0}).$$

Consequently

$$\det \begin{pmatrix} \log_2 2 & \log_2 M \\ \log_2 3 & \log_2 A \end{pmatrix} = -\log_2 M \log_2 3 \neq 0, \quad \det \begin{pmatrix} \log_3 2 & \log_3 M \\ \log_3 3 & \log_3 A \end{pmatrix} = \log_3 2 \log_3 A \neq 0.$$

After principal-unit normalization, the two-parameter orbit therefore has open closure in a two-dimensional p -adic torus at both structural primes. For a full-rank logarithm matrix B_p there is a constant c_p such that, for $z \neq 0$,

$$\left| \min_i v_p((B_p z)_i) - \min_i v_p(z_i) \right| \leq c_p.$$

Thus for $K \geq 2$, one fixed pair of distinct parameters in $\{0, \dots, K-1\}^2$ cannot have depths r_p at a finite set S of full-rank primes unless

$$\sum_{p \in S} r_p \log p \leq \log(K-1) + \sum_{p \in S} c_p \log p. \quad (338)$$

This is a genuine local packing ceiling, not a contradiction. The determinants of local logarithms are different transcendental periods at different places; they are not absolute values of one algebraic number and cannot be inserted directly into a product formula.

At a good prime $\ell \nmid 6MA$, after clearing denominators an integer relation $a \log_\ell 2 + b \log_\ell M = 0$ first implies $2^a M^b \in \ell^\mathbb{Z}$. Its ℓ -valuation is zero, hence it equals one; multiplicative independence then gives $a = b = 0$. This kernel step is essential before applying Baker–Brumer to a putative local shadow exponent. It is explicitly included here because omitting it would turn multiplicative independence into an invalid assertion about p -adic logarithms.

74.2 The interpolation denominator is a cokernel order

Let $K \geq 2$, $R = K^2$, $n = K - 1$, and enumerate the nodes and values

$$x_{ij} = 2^i 3^j, \quad y_{ij} = M^i A^j \quad (0 \leq i, j < K).$$

Let $E_K = (x_{ij}^r)_{(i,j), 0 \leq r < R}$ and let $Q_K \in \mathbb{Q}[X]$ be the unique interpolant with $Q_K(x_{ij}) = y_{ij}$. Define

$$\delta_K = \min\{d > 0 : dQ_K \in \mathbb{Z}[X]\}, \quad P_K = \delta_K Q_K,$$

and put

$$d_{ij}^{(K)} = \prod_{(a,b) \neq (i,j)} (x_{ij} - x_{ab}), \quad D_K^* = \text{lcm}_{i,j} |d_{ij}^{(K)}|.$$

The Lagrange basis and Smith normal form give three exact descriptions:

$$D_K^* = \min\{D > 0 : DE_K^{-1} \in M_R(\mathbb{Z})\} = s_R, \quad \delta_K = \text{ord}_{\mathbb{Z}^R/E_K \mathbb{Z}^R}[y], \quad (339)$$

where s_R is the largest invariant factor of E_K . If $UE_K V = \text{diag}(s_1, \dots, s_R)$, then more explicitly

$$\delta_K = \text{lcm}_{1 \leq r \leq R} \frac{s_r}{\gcd(s_r, (Uy)_r)}.$$

Thus all node-only denominator growth is separated from the order of the specialized value class. Any unexpectedly small denominator must be an arithmetic statement about that class, not merely a small Vandermonde determinant.

The exact endpoint valuation calculation finds the maximal dyadic barycentric denominator only in the top row and a maximizer at $(n, 0)$; triadically it lies only in the rightmost column with a maximizer at $(0, n)$. This yields two favorable, nonexhaustive branches:

$$M \text{ odd, } A \text{ even : } v_2([X^{R-1}]P_K) = 0, \quad (340)$$

$$3 \mid M, 3 \nmid A : v_3([X^{R-1}]P_K) = 0. \quad (341)$$

The restrictions in (341) are part of the theorem. They do not cover every hypothetical counterexample.

74.3 Exact scaling, roots, and an honest algebraic resultant

Let $F_K, G_K \in \mathbb{Z}[X]$ be the two degree- n residuals obtained from the dyadic and triadic boundary-dilation factorizations after multiplication by δ_K . In the first branch of (341), report 14 proves

$$\mathcal{F}_K(Z) = 2^{-n^2} F_K(2^n Z) \in \mathbb{Z}_2[Z], \quad \overline{\mathcal{F}_K}(Z) \doteq 1 - (Z - 1)^n, \quad (342)$$

with unit leading coefficient and $v_2(\mathcal{F}_K(0)) = v_2(M - 1)$. In the second branch,

$$\mathcal{G}_K(Z) = 3^{-n^2} G_K(3^n Z) \in \mathbb{Z}_3[Z], \quad \overline{\mathcal{G}_K}(Z) \doteq R_n(0) - \overline{A}R_n(Z), \quad R_n(Z) = \prod_{i=1}^n (Z - 2^i). \quad (343)$$

Consequently

$$v_2([X^r]F_K) \geq n(n - r), \quad v_3([X^r]G_K) \geq n(n - r),$$

and all roots have structural-prime depth at least n . The total excess depth is exactly $v_2(M - 1)$ or $v_3(A - 1)$, respectively. When n is odd the reductions are sufficiently separable to determine the exceptional root classes.

The categorical payoff is a single algebraic integer rather than unrelated local periods:

$$\mathcal{R}_{2,K} = \text{Res}(F_K(X), F_K(3X)), \quad (344)$$

$$\mathcal{R}_{3,K} = \text{Res}(G_K(X), G_K(2X)). \quad (345)$$

If the relevant resultant is nonzero, then

$$v_2(\mathcal{R}_{2,K}) \geq n^3 + n + v_2(M - 1), \quad (346)$$

$$v_3(\mathcal{R}_{3,K}) \geq n^3 + v_3(A - 1). \quad (347)$$

For odd n it is nonzero and equality holds. The finite algebraic hinge is the identity

$$P(3X) - P(X) = 2X \mathcal{D}_3(P), \quad \overline{\mathcal{D}_3(P)} = \overline{P}' \quad \text{in } \mathbb{F}_2[X], \quad (348)$$

where $\mathcal{D}_3(P) \in \mathbb{Z}[X]$ is an explicit geometric-sum quotient. Equation (348) is [\[kernel-verified Lean\]](#) in `DilationDifferenceQuotient`.

The resultants are the right kind of adelic objects, but their known archimedean heights are far larger than the local divisor in (347). No product-formula contradiction follows. The next missing theorem would have to control primitive resultant height or force compatibility between the two favorable branches; neither is supplied by local rank or by the Smith denominator alone.

74.4 Paired residuals and the full structural Smith profile

There is nevertheless more arithmetic in the case where the two favorable branches occur simultaneously:

$$M \text{ odd}, \quad A \text{ even}, \quad 3 \mid M, \quad 3 \nmid A. \quad (349)$$

This is a genuine restriction and is not known to follow from a hypothetical counterexample. Nor can a common change of exponent-grid directions manufacture it. Put

$$a = v_2(M), \quad c = v_2(A), \quad d = v_3(M), \quad b = v_3(A), \quad B_2 = \begin{pmatrix} 1 & 0 \\ a & c \end{pmatrix}, \quad B_3 = \begin{pmatrix} 0 & 1 \\ d & b \end{pmatrix}.$$

If a nonsingular common direction matrix C makes both B_2C and B_3C monomial—one nonzero in each row and column—then, and only then,

$$a = b = 0, \quad c, d \neq 0, \quad C \text{ itself is monomial.} \quad (350)$$

This is [\[kernel-verified Lean\]](#) as `simultaneous_cross_support_iff` in `GridTransformNoGo`. Translations do not change the direction matrices, and the theorem does not assume unimodularity; hence signed axis swaps, reflections, rectangular finite-index sublattices, and general $GL_2(\mathbb{Z})$ changes do not remove the branch restriction. A local shear may repair one structural prime but necessarily spoils the other. Dividing residuals by their contents is a different operation: it does not repair every fixed depth, and exact examples fail when the boundary degree equals one of the own-prime depths. There is, however, an exact eventual theorem after full content normalization; it is stated below with its strict threshold.

Write

$$F_K(X) = \sum_{r=0}^n f_r X^r, \quad G_K(X) = \sum_{r=0}^n g_r X^r, \quad \tau_r = \binom{n-r}{2}, \quad e_n = \binom{n}{2}.$$

The boundary commutator does not merely transfer the total structural content. Descending coefficient comparison gives the exact opposite-residual profiles

$$v_2(g_r) = \tau_r, \quad v_3(f_r) = \tau_r \quad (0 \leq r \leq n). \quad (351)$$

These complement the direct floors

$$v_2(f_r) \geq n(n-r), \quad v_3(g_r) \geq n(n-r)$$

from (342)–(343).

Here is the finite mechanism. In the notation of the boundary commutator, the coefficients of both L_3^- and L_3^+ have dyadic valuation τ_r : in the relevant elementary symmetric sum the unique minimum subset is $\{0, \dots, n-r-1\}$. The L_2^+ terms have the deep direct profile, while the cleared L_2^- term is strictly deeper. Compare the coefficient of X^{n+r} and descend from $r = n$ to $r = 0$. If an already known coefficient is shifted by $d \geq 1$, the doubled valuation gap is exactly

$$2(\tau_{r+d} + nd - \tau_r) = d(2r + d + 1) > 0. \quad (352)$$

The factor A raises the competing half by at least one, while M is a dyadic unit. Thus the coefficient equation forces $v_2(g_r) = \tau_r$. Interchanging $2, F, A, L_3$ with $3, G, M, L_2$ proves the triadic equality.

Let $S_K = \text{Sylv}(F_K, G_K)$ be the $2n \times 2n$ Sylvester matrix. Over the localization at either structural prime, its Smith-factor valuations are exactly

$$\underbrace{0, \dots, 0}_{n+1 \text{ times}}, \quad a_k = \frac{k(3k+1)}{2} \quad (1 \leq k < n). \quad (353)$$

Equivalently, if $\Delta_\ell(S_K)$ is the gcd of the ℓ -minors, then for $1 \leq m \leq n$

$$v_2(\Delta_{n+m}(S_K)) = v_3(\Delta_{n+m}(S_K)) = \mu_m := m \binom{m}{2} = \frac{m^2(m-1)}{2}. \quad (354)$$

In particular every $(n+m)$ -minor is divisible by 6^{μ_m} , and at the endpoint

$$\text{Res}(F_K, G_K) \neq 0, \quad v_2 \text{Res}(F_K, G_K) = v_3 \text{Res}(F_K, G_K) = \frac{n^2(n-1)}{2}. \quad (355)$$

For completeness, the Smith calculation is local and finite. Unit-normalize the direct residual to be monic and pivot on its shifted copies. This gives a unimodular equivalence

$$\text{Sylv}(f, g) \sim I_n \oplus T,$$

where T is multiplication by $g \bmod f$ on the monomial basis of $R[X]/(f)$. The direct coefficient floor ensures that reduction modulo f does not disturb the shallow coefficients. Below the diagonal the entry costs are $\binom{n-i+j}{2}$; above it the wrap terms are strictly deeper. For any m -matching of selected rows and columns, convexity and the injective-index bounds give total cost at least $m\binom{m}{2}$. The bottom- m by first- m minor has a unique term of exactly this cost. Successive differences

$$\mu_m - \mu_{m-1} = \frac{(m-1)(3m-2)}{2}$$

give (353). The convex matching inequality and the universal Toeplitz-minor divisibility are [\[kernel-verified Lean\]](#) in `SmithProfileCore`, whose endpoint is `pow_smithMu_dvd_weightedLowerToeplitz_minor`. The residual-specific commutator induction, wrap estimate, and exact-minor unit remain paper proofs.

The same profile in fact extends beyond the simultaneous favorable branch after dividing by the full integer contents. Assume only that M, A are positive integers and that the two cross ranks are nonzero,

$$a = v_2(M), \quad b = v_3(A), \quad c = v_2(A) > 0, \quad d = v_3(M) > 0, \quad n > \max(a, b).$$

For a nonzero $H \in \mathbb{Z}[X]$, put $\text{cont}(H) = \gcd_r |[X^r]H|$ and $H^{\text{pr}} = H / \text{cont}(H)$. If $c_K = [X^{K^2-1}]P_K$, the local contents are exactly

$$\begin{aligned} v_2(\text{cont } G_K) &= v_2(c_K), & v_2(\text{cont } F_K) &= v_2(c_K) + a, \\ v_3(\text{cont } F_K) &= v_3(c_K), & v_3(\text{cont } G_K) &= v_3(c_K) + b. \end{aligned} \quad (356)$$

Writing $F_K^{\text{pr}} = \sum f_r X^r$ and $G_K^{\text{pr}} = \sum g_r X^r$, one has

$$v_2(g_r) = v_3(f_r) = \binom{n-r}{2} \quad (0 \leq r \leq n), \quad (357)$$

$$v_2(f_{n-e}) \geq e(n-1) + 1, \quad v_3(g_{n-e}) \geq e(n-1) + 1 \quad (1 \leq e \leq n). \quad (358)$$

Consequently the primitive cross-Sylvester matrix has, at both 2 and 3, the full Smith profile (353) and determinantal-divisor orders (354). Thus content normalization enlarges the exact profile from $a = b = 0$ to every cross-rank pair once $n > \max(a, b)$; it does not force the cross-rank hypotheses $c, d > 0$ for a hypothetical counterexample.

The proof is coefficientwise. In the Lagrange expansion the relevant endpoint is uniquely minimal once $a < n$ (and symmetrically $b < n$). Quotient and divided-difference costs give (357), while the direct residual has the stronger descending floor (358). Pivoting the Sylvester matrix again gives $I_n \oplus T$. With row coordinate x and column coordinate j , a nonwrapping entry has doubled cost $(x+j+1)(x+j)$; every wrapping entry has a strictly larger lower bound. A compressed-permutation argument then gives $2m\binom{m}{2}$ for every m -matching, with a unique antidiagonal minimizer. The endpoint gap, direct floor, wrap inequalities, and finite assignment bound are [\[kernel-verified Lean\]](#) in `ContentNormalizedProfileFinite`; the Lagrange identities, valuation transfer, Sylvester reduction, and application to the actual residuals are the [\[paper proof in this report\]](#) paper

layer. The strict threshold is necessary: equality cases $(K, M, A) = (4, 24, 2)$ with $n = a = 3$ and $(3, 3, 18)$ with $n = b = 2$ have noncanonical primitive Smith profiles. Exact finite diagnostics found no failure above the threshold, but those calculations are unbundled evidence rather than part of the theorem.

The exact prime-power cancellation, node-divisor normalizations, change-of-variable formula for resultants, and the implication from the two normalized residue identities to (355) are [\[kernel-verified Lean\]](#) in `CrossResidualScaling`, `ResultantScaling`, `Report14NodeDivisorScaling`, and `Report14CrossResultantBridge`. The strongest kernel endpoint, `report14_actual_divisors_favorable_cross_resultant`, keeps the primary interpolant scalings and the two nonzero residue reductions as explicit hypotheses; the barycentric derivation of those hypotheses from (349) remains the paper layer.

The full resultant is not the only exact local object. For $1 \leq s \leq n$, let $T_s(F_K, G_K)$ be the leading $2s \times 2s$ Sylvester block and put

$$C_s(F_K, G_K) = \det T_s(F_K, G_K) = \text{psc}_{n-s}(F_K, G_K)$$

up to the conventional sign. The direct and opposite coefficient profiles give the exact principal-subresultant chain

$$C_s(F_K, G_K) \neq 0, \quad v_2(C_s(F_K, G_K)) = v_3(C_s(F_K, G_K)) = \mu_s = \frac{s^2(s-1)}{2}. \quad (359)$$

Indeed every nonzero determinant matching has total displacement s^2 . If E units of displacement are assigned to the direct rows, its doubled structural-prime cost is at least

$$2nE + \sum_i d_i(d_i - 1) \geq 2\mu_s,$$

with equality only when $E = 0$ and all s shallow-row displacements equal s . Thus the unique minimal term is $\pm f_n^s g_{n-s}^s$ at 2 and $\pm g_n^s f_{n-s}^s$ at 3. The sharp cost inequality, strict case, and equality classification are [\[kernel-verified Lean\]](#) in `DeepSubresultantCost`. Their application to the residual coefficient profiles is the paper layer.

Writing $c = [X^{K^2-1}]P_K$, the first column of T_s is divisible by c ; in the simultaneous favorable branch c is a unit at 2 and 3. Hence

$$\Theta_s = C_s/c \in \mathbb{Z} \setminus \{0\}, \quad v_2(\Theta_s) = v_3(\Theta_s) = \mu_s.$$

For $s = n - j$ with any fixed $j \geq 1$, this proper subresultant retains the full leading $n^3/2$ local mass. It still does not solve the height problem. Define the truncated top-coefficient norms by

$$H_{F,s}^2 = \sum_{d=0}^{\min(n, 2s-1)} |f_{n-d}|^2, \quad H_{G,s}^2 = \sum_{d=0}^{\min(n, 2s-1)} |g_{n-d}|^2.$$

The natural Hadamard bound satisfies

$$0 < |\Theta_s| \leq \frac{H_{F,s}^s H_{G,s}^s}{|c|}, \quad \frac{H_{F,s}^s H_{G,s}^s}{|c|} \geq 6^{s(K^2-2)},$$

so at $s \sim \rho n$ its deficit from the forced divisor is already

$$\left(\rho - \frac{\rho^3}{2}\right) (\log 6)n^3 + O(n^2) > 0.$$

Exact synthetic calculations likewise find a large primitive core: at $K = 4$, $(M, A) = (9, 2)$, and $s = 2$, removing $c6^{\mu_2}$ leaves a primitive leading coefficient with 93 decimal digits. This is an unbundled exact-computation diagnostic, not a common-exponent theorem. The deep chain is a genuine local strengthening, but a cancellation-sensitive primitive-height estimate would still be new.

The all-minor profile still does not close the argument. Structural zero minors prevent a literal application of the all-nonzero-minor inequality of Section 77.1. A support-aware Cauchy–Binet bound can be written as follows. Put $h_F = \|F_K\|_2$, $h_G = \|G_K\|_2$, $k = n + m$. Row normalization and the fact that every k -row subset of an invertible matrix has some nonzero k -minor give the necessary inequality

$$6^{2\mu_m} \sum_{r=m}^n \binom{n}{r} \binom{n}{n+m-r} h_F^{-2r} h_G^{-2(n+m-r)} \leq \binom{2n}{n+m}. \quad (360)$$

Even in the optimistic normalization where only the leading-coefficient costs remain, the order- n^3 logarithmic surplus at $m \sim \rho n$ is bounded by

$$\frac{\rho^3}{2} \log 6 - \log 2 - \rho \log 3 \leq -\log 2 < 0 \quad (0 \leq \rho \leq 1). \quad (361)$$

Thus no intermediate exterior level repairs the endpoint constant using coefficient norms and Hadamard/Cauchy–Binet alone. The displayed rate inequality is [\[kernel-verified Lean\]](#) in `CompoundEndpointRate`. A proof still needs a new primitive-height saving, a support-aware cancellation theorem stronger than these norm bounds, or a way to reach the simultaneous residue branch from every hypothetical counterexample.

The full interpolation denominator does not supply the missing saving either. There is an exact smallest-grid counterexample to any universal denominator-survival claim. For the synthetic favorable pair $(M, A) = (9, 2)$ and $n = 1$, the primitive interpolation numerator

$$P = 127X^3 - 1212X^2 + 3227X - 2082$$

takes the values $(60, 120, 540, 1080)$ at $(1, 3, 2, 6)$, so its interpolation denominator is $\delta = 60$. The two boundary residuals are

$$F = 5552 - 127X, \quad G = 3175X + 1041, \quad \text{Res}(F, G) = -17759807 = -13 \cdot 31 \cdot 127 \cdot 347.$$

Thus $\gcd(\delta, \text{Res}(F, G)) = 1$, and the rational residual resultant is the reduced fraction $-17759807/60^2$. The entire scalar factor introduced by clearing the two residual denominators can cancel. This certificate is [\[kernel-verified Lean\]](#) in `DenominatorCancellationCounterexample`. The pair $(9, 2)$ is test data, not a common-exponent pair; accordingly the example rules out a grid-denominator-only theorem but does not rule out a bound using the still-missing common- β arithmetic.

The obstruction persists asymptotically. For every $n \geq 1$, apply the square-grid construction of side $K = n + 1$ to the same fixed synthetic pair $(M, A) = (9, 2)$; write δ_n for the least interpolation denominator and F_n, G_n for its two boundary residuals. Put

$$A_2(n) = \sum_{j=1}^n v_2(3^j - 1), \quad A_3(n) = \sum_{i=1}^n v_3(2^i - 1).$$

The exact square-grid denominator formula and structural resultant order are

$$v_p(\delta_n) = \frac{n^3 + 2n^2 - n}{2} + A_p(n), \quad v_p \text{Res}(F_n, G_n) = \frac{n^2(n-1)}{2} \quad (p = 2, 3).$$

Since $A_p(n) \geq 0$, subtraction gives the [\[paper proof in this report\]](#) all- n inequality

$$2n v_p(\delta_n) - v_p \operatorname{Res}(F_n, G_n) \geq n^4 + \frac{3}{2}n^3 - \frac{1}{2}n^2.$$

Thus grid vanishing, collisions, and the local Smith profile alone require an $\Omega(n^4)$ cancellation tax before they can even reach the primitive resultant. This all- n comparison is a paper consequence of the displayed exact formulas; exact-integer checks through $n = 5$ serve only as diagnostics. Its synthetic scope is essential.

There is also an exact endpoint factorization, but it removes no certified structural-prime mass. Let $D = K^2 - 1$, let $Q_K \in \mathbb{Q}[M, A][X]$ be the generic rational interpolant on the square grid, and write $\lambda_K = [X^D]Q_K$, $\kappa_K = Q_K(0)$ in $\mathbb{Q}[M, A]$. Node inversion gives

$$\kappa_K(M, A) = (-1)^D \left(\prod_{i,j < K} 2^i 3^j \right) \lambda_K(M/2, A/3).$$

For the two generic rational boundary residuals U_K, V_K , every monomial of the padded Sylvester determinant contains at least one of their two top coefficients and at least one of their two constant coefficients: killing both tops lowers both degrees, while killing both constants gives the common root zero. Their endpoint formulas therefore give the [\[paper proof in this report\]](#) all- K divisibility identity

$$\operatorname{Res}_X(U_K, V_K) = \lambda_K \kappa_K H_K^{\operatorname{rat}}, \quad H_K^{\operatorname{rat}} \in \mathbb{Q}[M, A].$$

Thus the primitive numerator has the two universal Cramer endpoint factors, one annihilating the integer exponent curve at $0, \dots, D-1$ and the other at $1, \dots, D$. Exact symbolic factorization for $K = 2, 3, 4$ finds both factors with multiplicity one and a remaining core irreducible over $\mathbb{Q}[M, A]$ of bidegree $(1, 1), (6, 6), (15, 15)$, respectively; the three specialized cores at $(9, 2)$ are pairwise coprime and 2, 3-adic units. These finite calculations are retained as an unbundled exact-CAS diagnostic, not as [\[software computation\]](#) and not as an all- K irreducibility theorem. They give no stable primitive factor with which to offset the $\Theta(n^4)$ height scale; a deeper clustered subresultant would be a genuinely different object.

Status: Report 14 was audited independently from the common base. Structural-prime rank and Fermat-quotient statements overlapping report 11 are cross-referenced rather than double-counted. The integral dilation quotient and its mod-2 derivative reduction are [\[kernel-verified Lean\]](#) in `DilationDifferenceQuotient`. The paired-residual cancellation, resultant scaling, actual node-divisor normalization, conditional exact cross-resultant, and Toeplitz Smith-minor core are also [\[kernel-verified Lean\]](#) in the five modules named above. The exact denominator-cancellation certificate is [\[kernel-verified Lean\]](#) in `DenominatorCancellationCounterexample`; the common direction-change obstruction is [\[kernel-verified Lean\]](#) in `GridTransformNoGo`, and the sharp deep-subresultant cost core is [\[kernel-verified Lean\]](#) in `DeepSubresultantCost`. The finite inequalities behind the eventual content-normalized profile are [\[kernel-verified Lean\]](#) in `ContentNormalizedProfileFinite`, while their Lagrange, valuation, and residual application is the paper layer. The denominator-lattice, endpoint, scaled-root, opposite-coefficient, residual-specific full Smith, and self-dilation-resultant theorems are self-contained paper arguments, explicitly restricted to M odd and A even or to $3 \mid M$ and $3 \nmid A$ where stated. The report’s larger symbolic audit is retained only as an unbundled diagnostic: the displayed source does not reconstruct every claimed interpolation, Newton polygon, and resultant. No conclusion here proves the conjecture.

75 Fifteenth report: synchronized information and decoder complexity

Report 15 is independent of report 14 and of every earlier continuation. It treats a hypothetical semigroup as a finite information source. Several of its dynamical statements overlap the Sturmian probability and natural-boundary packages of reports 7 and 9; those are cross-referenced. Its genuinely new contribution is a sharp fixed-rank sum-channel calculation, an exact two-view decoder, quantitative degree obstructions for algebraic decoders, and a finite Fourier formulation of tied determinant cancellation.

75.1 One magnitude channel and a rank-two valuation profile

For $0 \leq n, k < L$ set

$$X_{n,k} = 2^n M^k = 2^{n+\beta k}, \quad Y_{n,k} = 3^n A^k = 3^{n+\beta k}.$$

The binary digit length of $X_{n,k}$ and ternary digit length of $Y_{n,k}$ are the same integer, $n + \lfloor \beta k \rfloor + 1$. Thus the two apparent size observations are one channel, not two independent constraints. Its entropy is $\log L + O_\beta(1)$, while the source has entropy $2 \log L$. Prefix-code redundancies are bounded quasiperiodic functions of the same rotation phase, so their rate vanishes.

The synchronized carry word has zero entropy rate and logarithmic block entropy; its exact prefix mean, variance, and two-shift covariance are already [\[kernel-verified Lean\]](#) in `SturmianProbability`, while the carry/floor generating-series identities are [\[kernel-verified Lean\]](#) in `SturmianGeneratingSeries`. Pólya–Carlson supplies a natural boundary only after the classical radius-one and non-eventual-periodicity hypotheses. That one-variable obstruction does not create a two-prime functional equation.

Fixed integer valuation observations have entropy growth equal to the rank of the displayed linear forms times $\log L + O(1)$. For the two output boxes the normalized profile is a rank-two representable polymatroid. Consequently every universal homogeneous linear rank or Shannon inequality is automatically satisfied at leading order. Any contradiction must see heights, growing moduli, or a subleading arithmetic term.

75.2 Fixed-rank sums are asymptotically unordered-pair codes

Let E be an N -element subset of a fixed finitely generated multiplicative group of positive algebraic numbers, of rank r . The Evertse–Schlickewei–Schmidt unit-equation theorem [\[39\]](#) implies that the nondegenerate solutions of $x_1 + x_2 = x_3 + x_4$ contribute only $O_r(N)$ beyond the two trivial permutations. Hence

$$\mathcal{E}_+(E) = 2N^2 - N + O_r(N), \tag{362}$$

$$|E + E| = \frac{N^2}{2} + O_r(N), \tag{363}$$

$$H_2(Z + Z') = 2 \log N - \log 2 + O_r(N^{-1}), \tag{364}$$

$$H(Z + Z') = 2 \log N - \log 2 + O_r((\log N)/N), \tag{365}$$

for independent uniform $Z, Z' \in E$. Thus the sum determines the unordered input pair with error $O_r(N^{-1})$. This is a strong exact-shape theorem, but it is compatible with the hypothetical rank-two boxes: their additive channel is already almost maximally informative.

75.3 Exact synchronized decoding and degree growth

For any real $p > 1$, strict convexity gives a finite injectivity statement with no number theory in its proof:

$$a + b = c + d, \quad a^p + b^p = c^p + d^p, \quad a, b, c, d > 0 \implies \{a, b\} = \{c, d\}. \quad (366)$$

Indeed, with $s = a + b$ the function $x \mapsto x^p + (s - x)^p$ is strictly decreasing on $[0, s/2]$. This complete unordered-pair decoder is [\[kernel-verified Lean\]](#) in `SynchronizedSumDecoder`; the formal theorem is `eq_or_swap_of_sum_eq_rpow_sum_eq`.

Exact decodability does not mean bounded algebraic complexity. Suppose $L \geq 2$ and $R = P/Q \in \mathbb{R}(z)$ has $\deg P, \deg Q \leq d$, no pole at $2^n + 1$, and

$$R(2^n + 1) = 3^n + 1 \quad (0 \leq n < L). \quad (367)$$

After cross multiplication and translation by one, the left side is a nonzero exponential polynomial whose bases lie in

$$\{2^j : 0 \leq j \leq d\} \cup \{3 \cdot 2^j : 0 \leq j \leq d\}.$$

These $2d + 2$ bases are distinct. The Chebyshev zero bound therefore gives

$$L \leq 2d + 1, \quad d \geq \left\lceil \frac{L - 1}{2} \right\rceil. \quad (368)$$

The cross-multiplied inequality $L \leq 2d + 1$ is [\[kernel-verified Lean\]](#) as `rational_sum_decoder_card_le` in `RationalSumDecoder`; the formal hypotheses are $\deg P, \deg Q \leq d$, $Q \neq 0$, and the cross-multiplied identities $P(2^n + 1) = (3^n + 1)Q(2^n + 1)$ for all $0 \leq n < L$. The rational-function formulation follows from its no-pole premise. More generally, a nonzero $F \in \mathbb{R}[z, w]$ of total degree at most d satisfying $F(2^n + 1, 3^n + 1) = 0$ for $0 \leq n < L$ gives

$$L \leq \binom{d + 2}{2} - 1 = \frac{d(d + 3)}{2}. \quad (369)$$

The scope is the displayed exponential slice. A completion criterion may use these bounds only when its synchronized family contains $(2^n, 1)$ for arbitrarily long initial ranges; they do not prove degree growth on every unbounded collection of semigroup pairs.

75.4 Fourier separation identifies, but does not solve, tied cancellation

Let $B \in \{\pm 1\}$ be a balanced parity bit and Z a random element of a finite abelian group G . The signed measure

$$\nu(g) = \mathbb{E}[B \mathbf{1}_{Z=g}]$$

has Fourier coefficients $\hat{\nu}(\chi) = \mathbb{E}[B \overline{\chi(Z)}]$. Finite Fourier inversion and Parseval give

$$B \text{ independent of } Z \iff \hat{\nu}(\chi) = 0 \quad (\chi \in \widehat{G}), \quad (370)$$

$$\text{TV}(\mathcal{L}(B, Z), \mathcal{L}(B) \otimes \mathcal{L}(Z)) \leq \frac{1}{2} \left(\sum_{\chi \in \widehat{G}} |\hat{\nu}(\chi)|^2 \right)^{1/2}. \quad (371)$$

If B is determined by Z , at least one character correlation has magnitude $|G|^{-1/2}$. Applied to tied determinant terms, this says exactly what would force noncancellation: a legal twist family spanning all residue characters while preserving the archimedean estimate and the tied fiber. No current determinant supplies that family, so (371) is a precise target rather than a completed argument.

Status: Report 15 was audited independently from the common base. Its exact ordinary/power-sum decoder is [\[kernel-verified Lean\]](#) in `SynchronizedSumDecoder`. The rational decoder degree bound is [\[kernel-verified Lean\]](#) in `RationalSumDecoder`. The algebraic decoder degree bound and finite Fourier lemma are unconditional paper proofs; the fixed-rank entropy estimates use the classical unit-equation theorem. Sturmian and Pólya–Carlson material overlapping reports 7 and 9 is not double-counted. The report’s uniform-decoder criterion is stated only for families containing the exponential slice $(2^n, 1)$ through arbitrarily large initial ranges. Character-twisted determinant separation remains an open construction. None of the entropy identities by itself proves the conjecture.

76 Sixteenth report: transfer spectra and mixed Mahler resonance

Report 16 is another independent sibling of reports 1–15: it starts from the same unified base and asks whether classical special functions can amplify the multiplicative rank of the two known outputs. Its generic spectrum theorem strengthens the existing $(2, 3)$ classification, while its gamma, Hurwitz, polylogarithmic, Barnes, double-sine, and q -digamma packages diagnose several exact but rank-preserving reformulations. The report’s recorded metadata searches all failed, so no “no proof found” conclusion is imported here; only mathematical claims with an internal proof or an identified classical input are used.

76.1 The algebraic-base spectrum

Let $a, b, c > 0$ be algebraic, assume that the nonnegative monomials in a, b are unique, and let $x \notin \mathbb{Q}$ with a^x, b^x algebraic. Write

$$\text{Rel}_{a,b}(c) \iff \exists n > 0 \exists i, j, k, l \in \mathbb{N}_0 : a^k b^l c^n = a^i b^j. \quad (372)$$

Then

$$\boxed{c^x \in \overline{\mathbb{Q}} \iff \text{Rel}_{a,b}(c).} \quad (373)$$

Indeed, the forward implication follows from the algebraic Six Exponentials theorem unless the three-base nonnegative monomial map has a collision; the reverse implication follows by raising (372) to x and extracting a positive real n th root. This is [\[kernel-verified Lean\]](#) as `AlgebraicBaseSpectrum.real_rpow_isAlgebraic_iff_hasPositivePairPowerRelation`; the exact collision equivalence is `AlgebraicBaseSpectrum.tripleMonomial_injective_iff`.

For a hypothetical nonintegral solution and $(a, b) = (2, 3)$, this recovers the exact statement that the positive algebraic bases with algebraic x th power are precisely the divisible hull of $\langle 2, 3 \rangle$. Thus no positive algebraic third base can simultaneously have an algebraic output and supply an independent logarithmic direction. Gamma quotients, radicals, norms, and structural-prime constructions do not evade this theorem once their positivity and algebraicity are established.

76.2 Entire cancellation and rank-preserving transfer formulas

For an integer $b \geq 2$, the apparent meromorphic quotient

$$Q_b(z) := \frac{\Gamma(-z)}{\Gamma(1-b^z)} \quad (374)$$

extends to an entire function: at every $k \in \mathbb{N}_0$ the two simple poles cancel with a nonzero ratio. Its other zeros are the nonintegral points for which b^z is a positive integer. Since the two vertical phase lattices for bases 2 and 3 meet only on the real axis,

$$\text{Alaoglu–Erdős} \iff Z(Q_2) \cap Z(Q_3) = \emptyset. \quad (375)$$

At a nonintegral real zero with $b^x = B \in \mathbb{N}$ one has

$$Q'_b(x) = \Gamma(-x)(-1)^B B! \log b. \quad (376)$$

These paper identities remove all known integral solutions exactly, but they do not create a two-dimensional common zero lattice; a common Barnes divisor would therefore be spurious.

Gauss multiplication, Hurwitz distribution, and the Cartier identity for $F_x(z) = \text{Li}_{-x}(z)$ all yield the same multiplier q^x . Under a counterexample their normalized values are algebraic exactly for 3-smooth q , by (373). The dichotomy

$$\text{Li}_{-x}(z) \text{ is D-finite over } \mathbb{C}(z) \iff x \in \mathbb{N}_0 \quad (377)$$

is classical: nonintegral order produces infinitely many independent fractional-power monodromy germs. Likewise, after defining $a_0 = 0$, the sequence $a_n = n^x$ is q -regular exactly for integral x . Integer Cartier eigenvalues at 2 and 3, however, do not imply either finite holonomic rank or finite q -kernel rank; that missing implication is another form of the original wall.

76.3 A finite Mahler identity and the corrected resonance theorem

For a commutative ring and finite cutoffs define

$$H_{A,B}(z) := \sum_{0 \leq a < A} \sum_{0 \leq b < B} M^a N^b z^{2^a 3^b}. \quad (378)$$

The exact rectangular inclusion–exclusion identity

$$H_{A+1,B+1}(z) = z + MH_{A,B+1}(z^2) + NH_{A+1,B}(z^3) \quad (379)$$

$$- MNH_{A,B}(z^6) \quad (380)$$

is [\[kernel-verified Lean\]](#) as `mixedMahlerTrunc_succ_succ` in `MixedMahlerFinite`. It is a finite algebra theorem only; convergence, natural boundaries, Mellin transforms, and asymptotics are not hidden in its statement.

For positive integers M, N , the limiting integer-coefficient series

$$H_{M,N}(z) = \sum_{a,b \geq 0} M^a N^b z^{2^a 3^b} \quad (381)$$

has radius one and a non-eventually-periodic zero pattern, so the classical Pólya–Carlson/Skolem–Mahler–Lech argument [40, 41] makes the unit circle a natural boundary. This happens for both integral and hypothetical nonintegral exponents and is therefore only a diagnostic.

For $M, N \geq 2$, put

$$\alpha = \frac{\log M}{\log 2}, \quad \beta = \frac{\log N}{\log 3}. \quad (382)$$

Termwise Mellin integration initially gives

$$\int_0^\infty H_{M,N}(e^{-t})t^{s-1} dt = \frac{\Gamma(s)}{(1 - M2^{-s})(1 - N3^{-s})} \quad (\Re s > \max\{\alpha, \beta\}). \quad (383)$$

The two geometric pole lattices meet exactly when $\alpha = \beta$, and then only at the real point $s = x$. Consequently resonance gives

$$H_{M,N}(e^{-t}) = t^{-x} \left[\frac{\Gamma(x)}{\log 2 \log 3} \log \frac{1}{t} + B + P_2(\log(1/t)) + P_3(\log(1/t)) \right] + O(t^{-x+\eta}), \quad (384)$$

where P_2, P_3 are absolutely convergent analytic periodic functions with periods $\log 2, \log 3$. The restriction $M, N \geq 2$ is essential: $M = 1$ or $N = 1$ introduces an unrelated Gamma-pole collision at $s = 0$. The Barnes kernel

$$\frac{1}{(1 - e^{-v \log 2})(1 - e^{-v \log 3})} = \sum_{a,b \geq 0} e^{-v(a \log 2 + b \log 3)} \quad (385)$$

is valid as a geometric series for $\Re v > 0$; the corresponding Barnes double-zeta integral starts in $\Re w > 2$, $\Re a > 0$ and is then continued. The double pole detects equality of the two exponents, not their integrality.

76.4 Double-sine, q -digamma, and the surviving target

The reciprocal double-sine normalization with periods $\log 2, \log 3$ packages the two outputs as adjacent shift quotients at $z_x = ix \log 2 \log 3 / (2\pi)$. The shift compatibility is built into the function and supplies no contradiction. Separately, Borwein’s Padé theorem gives the unconditional irrationality of

$$L_b(B) = \sum_{k \geq 0} \frac{1}{Bb^k - 1} = -\frac{\psi_{1/b}(x) + \log(1 - b^{-1})}{\log b} \quad (B = b^x \in \mathbb{N}, b \geq 2). \quad (386)$$

This one-base irrationality [42] also holds for every genuine integral exponent. Finally, ${}_1F_0(-x; -; -z) = (1 + z)^x$ (continued analytically at the displayed arguments) makes the rank-amplification target explicit: the known values at $1 + z = 2, 3$ do not algebraize the value at $1 + z = 5$. A cross-base value theorem doing so would contradict (373), but no such theorem is presently available.

Status: Report 16 was audited independently from the common base. The generic algebraic-base spectrum is [\[kernel-verified Lean\]](#) in `AlgebraicBaseSpectrum`; the finite mixed Mahler identity is [\[kernel-verified Lean\]](#) in `MixedMahlerFinite`. The entire quotient, corrected $M, N \geq 2$ Mellin resonance, and transfer identities are self-contained paper arguments. The holonomic dichotomy, Pólya–Carlson boundary, Barnes continuation, and Borwein irrationality use named classical inputs. Double-sine rigidity, Cartier-to-holonomy, and all proposed rank-amplification statements remain open. Nothing in the section proves the conjecture.

77 Seventeenth report: weighted Smith profiles and rational Loewner clusters

Report 17 is independent of reports 1–16. It replaces the single determinant by the full sequence of determinantal divisors and compares that arithmetic profile with every exterior singular-value level. It also develops a rational cross-Loewner ensemble and proves that ordinary probability and entropy are already saturated by the hypothetical semigroup. The new certificates are exact sufficient criteria, not verified contradictions. The basic exterior/Frobenius identity overlaps report 13, the fractional Loewner integral and rational positivity overlap report 10, and source-recovery entropy overlaps report 15. The genuinely new refinements are the weighted all-minor inequality, the p^t -level modular divisor, the full geometric-Vandermonde Smith form, and the level-by-level Loewner singular hierarchy.

77.1 From one determinant to every exterior level

For a full-rank $A \in \mathbb{Z}^{n \times n}$ let $\delta_k(A)$ be the gcd of its k -minors and let $\sigma_1 \geq \cdots \geq \sigma_n$ be its singular values. Cauchy–Binet in exterior form gives

$$\sum_{|I|=|J|=k} \det(A_{I,J})^2 = e_k(\sigma_1^2, \dots, \sigma_n^2). \quad (387)$$

If all k -minors are nonzero, then

$$e_k(\sigma(A)^2) \geq \binom{n}{k}^2 \delta_k(A)^2. \quad (388)$$

Thus, if $A = B + R$, $\text{rank } B \leq r < k$, $\|R\|_2 \leq \varepsilon$, and $\|A\|_2 \leq S$, then

$$\boxed{\sqrt{\binom{n}{k}} \delta_k(A) \leq S^r \varepsilon^{k-r}.} \quad (389)$$

Every intermediate k is a possible contradiction level.

For a diagonal gauge $K = D^{-1}AE^{-1}$ with positive diagonal entries ρ_i, κ_j , the exact weighted identity is

$$e_k(\sigma(K)^2) = \sum_{I,J} \frac{\det(A_{I,J})^2}{\rho_I^2 \kappa_J^2} \geq \delta_k(A)^2 e_k(\rho^{-2}) e_k(\kappa^{-2}). \quad (390)$$

This retains Smith mass after removing analytically irrelevant row and column scales, unlike a determinant-only normalization.

77.2 Multilevel residue collapse and a geometric Smith model

For a d -variable integer monomial-evaluation matrix, let $0 \leq \nu_{p,1} \leq \cdots \leq \nu_{p,n}$ be the valuations of its Smith factors and put $r_{p,t} = \#\{i : \nu_{p,i} < t\}$. The exact Ferrers-level identity and residue-class bound are

$$v_p(\delta_k) = \sum_{t \geq 1} (k - r_{p,t})_+, \quad r_{p,t} \leq p^{td}. \quad (391)$$

Consequently every such matrix satisfies

$$D_{k,d} := \prod_p p^{\sum_{t \geq 1} (k - p^{td})_+} \mid \delta_k. \quad (392)$$

For fixed d ,

$$\log D_{k,d} = \left(\frac{d}{d+1} + o(1) \right) k^{1+1/d}; \quad (393)$$

the bivariate constant is $2/3$. This paper theorem applies to arbitrary supports, but its universal size is lower order in the natural global determinant budgets.

The one-variable geometric Vandermonde

$$V_n(q) = (q^{ij})_{0 \leq i,j < n}, \quad q \geq 2, \quad (394)$$

has the exact Smith factors

$$s_i(q) = q^{\binom{i}{2}} \prod_{r=1}^i (q^r - 1), \quad 0 \leq i < n. \quad (395)$$

The proof changes from monomials to the Newton basis $\prod_{r < i} (X - q^r)$, leaving a diagonal matrix times an upper unitriangular Gaussian-binomial matrix. Hence

$$\log \delta_k(V_n(q)) = \frac{\log q}{3} k^3 + O_q(k^2). \quad (396)$$

This large one-base Smith mass occurs for every ordinary integer q and is therefore a calibration and no-go, not evidence of exceptional two-base structure. The complete determinant-one two-sided reduction, the integral q -Pascal identity, the divisibility chain, and positivity for $q \geq 2$ are [\[kernel-verified Lean\]](#) in `GeometricVandermondeSmith`; the formal endpoint is `geometricVandermonde_isUnimodularEquivalent_qSmithDiagonal`.

77.3 Rational cross-Loewner clusters

Write a hypothetical exponent as $x = m + \alpha$ with $0 < \alpha < 1$. For u, v in the dense group $\{2^a 3^b : a, b \in \mathbb{Z}\}$, both u^α, v^α are rational, so the cross-Loewner kernel

$$L_\alpha(u, v) = \frac{u^\alpha - v^\alpha}{u - v} = \frac{\sin(\pi\alpha)}{\pi} \int_0^\infty \frac{t^\alpha}{(u+t)(v+t)} dt \quad (397)$$

is rational at distinct group nodes. Andreief's identity and the Cauchy determinant give, for disjoint increasing node sets,

$$\det(L_\alpha(u_i, v_j)) = \frac{c_\alpha^n}{n!} \Delta(u) \Delta(v) \quad (398)$$

$$\times \int_{(0,\infty)^n} \frac{\prod t_\ell^\alpha \Delta(t)^2}{\prod_{i,\ell} (u_i + t_\ell)(v_i + t_\ell)} dt, \quad (399)$$

which is strictly positive. For affine clusters $u_i = 1 + \varepsilon \xi_i$, $v_j = 1 + \varepsilon \eta_j$, Selberg's integral yields

$$\det(L_\alpha(u_i, v_j)) = \kappa_{n,\alpha} \Delta(\xi) \Delta(\eta) \varepsilon^{n(n-1)} (1 + O(\varepsilon)), \quad (400)$$

with an explicit positive gamma-product constant. Orthogonal projection in the Stieltjes feature space gives the sharper individual hierarchy

$$\sigma_j = \Theta(\varepsilon^{2(j-1)}) \quad (1 \leq j \leq n) \quad (401)$$

for fixed nondegenerate affine clusters; the determinant constant is the corresponding Selberg integral [46].

If $A = DKE$ clears a rational cluster to integers, combining (390) and (401) forces

$$\delta_k(A) \left(\frac{e_k(D^{-2})e_k(E^{-2})}{\binom{n}{k}} \right)^{1/2} \leq C_{n,k,\alpha} \varepsilon^{k(k-1)}. \quad (402)$$

A finite violation would prove the conjecture. Naive numerator/denominator clearing is exponentially more expensive than this polynomial target; reduced Smith content or genuine row–column cancellation is still missing.

77.4 Positivity and entropy no-go theorems

The most natural 3-smooth Gram matrices are infinitely divisible: if

$$K_{ij} = 2^{\langle a_i, a_j \rangle} 3^{\langle b_i, b_j \rangle}, \quad (403)$$

then $K^{ot} \succeq 0$ for every $t \geq 0$, and under a counterexample K^{ox} is integral. Thus positivity of Hadamard powers cannot distinguish integral from nonintegral x without an additional arithmetic positivity principle.

Probability is equally tight. If N, K are independent uniform variables on $\{0, \dots, H\}$ and

$$X_H = 2^N M^K, \quad Y_H = 3^N A^K, \quad (404)$$

then one external-prime valuation and the structural valuation recover (N, K) from either output separately. Hence

$$\mathbf{H}(X_H) = \mathbf{H}(Y_H) = \mathbf{H}(X_H, Y_H) = 2 \log(H+1), \quad I(X_H; Y_H) = 2 \log(H+1). \quad (405)$$

The hypothetical semigroup saturates, rather than violates, ordinary anti-concentration and Shannon inequalities. Random minors and modular scouting remain useful discovery tools only when followed by an exact Smith or interval certificate.

Status: Report 17 was audited independently from the common base. The exact geometric-Vandermonde Smith form is [\[kernel-verified Lean\]](#) in `GeometricVandermondeSmith`. The exterior-minor identities, weighted Smith inequalities, multilevel modular divisor, Loewner–Andreief/Selberg formulas, singular hierarchy, and entropy saturation remain self-contained paper arguments with the stated classical linear-algebra and integral inputs. The universal modular divisor is a baseline, not a closing estimate; the geometric Smith form is a one-base no-go; and the Loewner certificate still requires reduced denominator or Smith information not supplied by height clearing. None of these results proves the conjecture.

78 Eighteenth report: logarithmic spectral rank and approximation capacity

Report 18 is an independent asymptotic reading of the common exact conditional structure. It regards the candidate set as the spectrum of a diagonal operator, resolves the associated heat trace by Mellin inversion, records a second mixed support equation, and compares the best polynomial approximation rate on a fixed multiplicative window with the number of conditional integer nodes. These constructions distinguish a formal rank-one semigroup from a formal rank-two semigroup; by themselves they do not use the additional arithmetic fact that the synchronized ϑ th powers are also integers.

78.1 Spectral zeta and logarithmic Weyl rank

On $\ell^2(\mathcal{S})$ define $De_x = 2^x e_x$ and $Ee_x = 3^x e_x$. Positive functional calculus gives $E = D^\vartheta$. For $\Re s > 0$, the exact conditional structure gives

$$Z_D(s) = \text{Tr}(D^{-s}) = \begin{cases} (1 - 2^{-s})^{-1}, & \mathcal{S} = \mathbb{N}_0, \\ ((1 - 2^{-s})(1 - M^{-s}))^{-1}, & \mathcal{S} = \mathbb{N}_0 + \mathbb{N}_0\beta, \end{cases} \quad (406)$$

and $Z_E(s) = Z_D(\vartheta s)$. Thus the origin is a simple pole in rank one and a double pole in rank two. Equivalently, for

$$N_D(\Lambda) = \#\{x \in \mathcal{S} : 2^x \leq \Lambda\}, \quad (407)$$

the exact conditional count, together with the paper-level Baker–discrepancy refinement, gives

$$N_D(\Lambda) = \frac{\log \Lambda}{\log 2} + O(1) \quad \text{in rank one,} \quad (408)$$

$$N_D(\Lambda) = \frac{(\log \Lambda)^2}{2 \log 2 \log M} + \left(\frac{1}{2 \log 2} + \frac{1}{2 \log M} \right) \log \Lambda + O_M((\log \Lambda)^{1-\delta}) \quad \text{in rank two.} \quad (409)$$

Consequently the limit

$$\dim_{\log} D := \lim_{\Lambda \rightarrow \infty} \frac{\log N_D(\Lambda)}{\log \log \Lambda} \quad (410)$$

is respectively 1 or 2. This is an exact reformulation of the monoid dichotomy, not an independent method for deciding which case holds.

78.2 Complete heat-trace expansion

For a formal independent pair $(2, M)$ put $a = \log 2$, $b = \log M$ and

$$H_{a,b}(t) = \sum_{n,k \geq 0} \exp(-te^{an+bk}). \quad (411)$$

Mellin transformation gives

$$\int_0^\infty H_{a,b}(t) t^{s-1} dt = \frac{\Gamma(s)}{(1 - e^{-as})(1 - e^{-bs})} \quad (\Re s > 0). \quad (412)$$

Since 2 and M are multiplicatively independent, Baker’s theorem supplies polynomial lower bounds for the near-collisions of the two imaginary pole lattices. The exponential vertical decay of Γ then makes the residue series absolutely convergent. With $L = \log(1/t)$, contour shifting yields, for every $R \geq 0$,

$$H_{a,b}(t) = Q_{a,b}(L) + P_a(L) + P_b(L) \quad (413)$$

$$+ \sum_{r=1}^R \frac{(-1)^r t^r}{r!(1 - e^{ar})(1 - e^{br})} + O_{a,b,R}(t^{R+1/2}), \quad (414)$$

where

$$Q_{a,b}(L) = \frac{L^2}{2ab} + \left(\frac{1}{2a} + \frac{1}{2b} - \frac{\gamma}{ab} \right) L + C_{a,b}, \quad (415)$$

$$P_a(L) = \sum_{\ell \neq 0} \frac{\Gamma(2\pi i \ell / a)}{a(1 - e^{-2\pi i \ell b / a})} e^{2\pi i \ell L / a}, \quad (416)$$

and P_b is the symmetric series. The two nonzero frequency sets are disjoint; hence the bounded remainder contains a genuine quasiperiodic component. In rank one the leading term is L/a with one periodic frequency lattice. These paper formulas are mathematically sharper than the rough count, but the output heat trace is only the rescaling $Z_E(s) = Z_D(\vartheta s)$ and supplies no second independent constraint.

78.3 The finite two-scale support equation

For natural scales p, q and a commutative ring, set

$$F_{A,B}(z) = \sum_{0 \leq a < A} \sum_{0 \leq b < B} z^{p^a q^b}. \quad (417)$$

The exact finite identity

$$F_{A+1,B+1}(z) = z + F_{A,B+1}(z^p) + F_{A+1,B}(z^q) - F_{A,B}(z^{pq}) \quad (418)$$

is [\[kernel-verified Lean\]](#) as `twoScaleSupportTrunc_succ_succ` in `TwoScaleSupportFinite`. It is coefficientwise finite and has no analytic hypotheses.

For multiplicatively independent integers 2, M , passage to the convergent support series

$$F_{2,M}(z) = \sum_{n,k \geq 0} z^{2^n M^k}, \quad |z| < 1, \quad (419)$$

gives

$$F_{2,M}(z) - F_{2,M}(z^2) - F_{2,M}(z^M) + F_{2,M}(z^{2M}) = z. \quad (420)$$

Its 0–1 coefficient sequence has density-zero infinite support and is not eventually periodic, so Pólya–Carlson gives the unit circle as a natural boundary. Rank-one lacunary series already have the same boundary; only the radial singularity order sees the rank. Existing Mahler theorems for two separate independent functional equations do not apply to this one mixed inclusion–exclusion equation.

78.4 A sharp fixed-window node-budget barrier

Fix $h > 0$. On $[1, e^h]$ the nearest singularity of u^ϑ maps, under the affine normalization to $[-1, 1]$, to the Bernstein ellipse parameter

$$\rho(h) = \coth(h/4). \quad (421)$$

Therefore the least polynomial degree needed to approximate x^ϑ uniformly within $1/2$ on $[X, e^h X]$ obeys

$$d_{\min}(X, h) \geq \left(\frac{\vartheta}{\log \rho(h)} + o(1) \right) \log X. \quad (422)$$

Under a counterexample the number of conditional nodes in the same window is

$$K_X(h) = \frac{h}{\beta(\log 2)^2} \log X + o(\log X). \quad (423)$$

The elementary maximization

$$\sup_{h>0} h \log \coth(h/4) = 2 \operatorname{arsinh}(1)^2 = 1.5536387997 \dots \quad (424)$$

occurs at $h = 2 \operatorname{arsinh}(1)$. Using the kernel-verified lower bound $\beta > 12$ therefore gives

$$\limsup_{X \rightarrow \infty} \frac{K_X(h)}{d_{\min}(X, h)} \leq \frac{h \log \coth(h/4)}{\beta \log 2 \log 3} < 0.170020. \quad (425)$$

Thus a fixed-window strategy that first approximates by a degree- d polynomial to sub-unit accuracy and then demands at least $d + 2$ integer nodes is starved by a factor greater than 5.88, even before coefficient height and denominator taxes. This does not exclude rational, multivariate, semigroup-adapted, or growing-window approximants.

The report also compares the conditional $(\log X)^2$ set with recent irrational-power additive-relation bounds. Unequal-length relations such as $(2m)^\vartheta = 3m^\vartheta$ already occur at the scale allowed by those theorems, while no extra equal-length energy is forced. Because the relevant preprints and constants are time-sensitive, this synthesis retains only the mathematical threshold diagnosis and does not elevate a dated literature search to a theorem.

Status: Report 18 was audited independently from the common base. The finite two-scale support identity is [\[kernel-verified Lean\]](#) in `TwoScaleSupportFinite`. The zeta/Weyl reformulations use the kernel-verified exact monoid and leading counting asymptotic, with the displayed power-saving remainder kept at paper level. The complete heat expansion and sharp fixed-window capacity theorem are paper arguments using, respectively, classical Mellin–Baker input and classical best-polynomial-approximation asymptotics. The natural boundary is classical after the finite support algebra. None of the spectral, heat, Mahler, approximation, or additive signatures eliminates a formal rank-two integer semigroup; the missing step is still the arithmetic coupling $M^\vartheta \in \mathbb{N}$.

79 Nineteenth report: arithmetic approximation and geometric Newton interpolation

Report 19 is an independent approximation-theoretic sibling of the preceding reports. Its fixed-window optimization converges on the same capacity constant as report 18, but starts from the

integer-valued-polynomial lattice on the 3-smooth nodes rather than from a spectral count. Its genuinely complementary finite content is an exact interpolation certificate, the cofactor formula for discrete minimax error, and a geometric Newton expansion with an explicit first extrapolation error. These results isolate the arithmetic unit barrier; they do not manufacture the missing integer-valued approximant.

79.1 Subunit locking and the optimal multiplicative block

Put $\mathcal{B} = \{2^a 3^b : a, b \in \mathbb{N}_0\}$. Under a hypothetical solution $\beta, q^\beta \in \mathbb{Z}$ for every $q \in \mathcal{B}$. If $E \subset \mathcal{B}$ is finite and a degree- d polynomial P is integer-valued on E , then

$$\max_{q \in E} |q^\beta - P(q)| < 1 \implies q^\beta = P(q) \quad (q \in E). \quad (426)$$

The generalized power system $1, t, \dots, t^d, t^\beta$ is Chebyshev on $(0, \infty)$, so a nonintegral β makes (426) impossible as soon as $|E| \geq d + 2$. The same argument gives an exact unit floor for every integer-coefficient Müntz polynomial with distinct exponents in $\mathbb{N}_0 + \mathbb{N}_0\beta$.

For $c > 1$, the 3-smooth nodes in $[X, cX]$ have cardinality

$$\frac{\log c}{\log 2 \log 3} \log X + o(\log X), \quad (427)$$

whereas Bernstein approximation to t^β on that block requires degree at least $\beta \log X / \log \rho(c) + o(\log X)$, with

$$\rho(c) = \frac{\sqrt{c} + 1}{\sqrt{c} - 1}. \quad (428)$$

The exact optimization is

$$\max_{c > 1} \log c \log \rho(c) = 2 \log^2(1 + \sqrt{2}), \quad c_* = 3 + 2\sqrt{2}. \quad (429)$$

Consequently this particular method could have a positive degree-versus-node margin only for

$$M = 2^\beta < \exp\left(\frac{2 \log^2(1 + \sqrt{2})}{\log 3}\right) = 4.113124586 \dots \quad (430)$$

Among nonpowers of two only $M = 3$ lies in that formal range, and the accepted finite corpus already excludes it (indeed it proves $M \geq 4096$ for a counterexample). Thus the exact constant is a sharp no-go for ordinary fixed-block polynomial approximation, consistent with the report-18 capacity calculation.

79.2 Dyadic certificates and the discrete minimax unit barrier

Let $0 \leq n_0 < \dots < n_d$, and let $I_{\mathbf{n}} \in \mathbb{Q}[T]$ be the degree-at-most- d interpolant with $I_{\mathbf{n}}(2^{n_j}) = 3^{n_j}$. For an integer M outside the node set, write $I_{\mathbf{n}}(M) = a/b$ in lowest terms, $b > 0$. The exact certificate is

$$\left| M^\vartheta - \frac{a}{b} \right| < \frac{1}{b} \implies M^\vartheta \notin \mathbb{Z}. \quad (431)$$

If the target were integral, its nonzero difference from a/b would have size at least $1/b$; equality with the interpolant is excluded by the $d + 2$ -zero Chebyshev bound. The quadratic interpolants

$$I_{0,1,2}(T) = \frac{T^2 + 3T - 1}{3}, \quad I_{1,2,3}(T) = \frac{T^2 + 6T - 4}{4} \quad (432)$$

give short exact exclusions for $M = 3, 5, 6, 7$. They are conceptually transparent certificates, not an improvement on the much larger kernel-verified finite range.

More generally, let A be an $(N + 1) \times N$ generalized evaluation matrix, f the target column, and C_i the signed maximal cofactors of A . With

$$D = \det[A \ f], \quad S = \sum_{i=0}^N |C_i|, \quad (433)$$

strict total positivity gives the exact real minimax value

$$\min_{c \in \mathbb{R}^N} \|f - Ac\|_\infty = \frac{|D|}{S}. \quad (434)$$

When A and f are integral, clearing this canonical denominator produces $z \in \mathbb{Z}^N$ with

$$Sf - Az = D (\operatorname{sgn} C_i)_{i=0}^N. \quad (435)$$

Thus the small real residual D/S becomes the nonzero integer residual D . Smith content or a different integer-valued basis must save part of S ; ordinary coefficient clearing cannot cross the unit threshold.

79.3 The geometric Newton shadow

For integers $q \geq 2$, $r \geq 1$, define

$$B_k^{(q)}(X) = \prod_{i=0}^{k-1} \frac{X - q^i}{q^k - q^i}, \quad a_k(q, r) = \prod_{i=0}^{k-1} (r - q^i). \quad (436)$$

At $X = q^n$, the first factor is the Gaussian binomial $\begin{bmatrix} n \\ k \end{bmatrix}_q \in \mathbb{Z}$, and the exact Newton identity is

$$r^n = \sum_{k=0}^n \begin{bmatrix} n \\ k \end{bmatrix}_q \prod_{i=0}^{k-1} (r - q^i). \quad (437)$$

Hence the truncation through $k = d$ is integer-valued at every geometric node, interpolates r^n for $n \leq d$, and has first omitted error

$$r^{d+1} - P_d^{q,r}(q^{d+1}) = \prod_{i=0}^d (r - q^i). \quad (438)$$

Unless r is a nonnegative integral power of q , the Newton series never terminates and this extrapolation product grows rather than shrinks. This complements the exact geometric Smith form of report 17: one-axis integer-valued interpolation is abundant, but it supplies no cross-axis smallness.

Status: Report 19 was audited independently from the common base. The unit-locking, optimal-block, dyadic-certificate, minimax-cofactor, and arithmetic-alternant statements are self-contained paper arguments using the stated classical Chebyshev/Bernstein inputs. The fixed-window constant overlaps report 18 and is already incompatible with the kernel-verified counterexample range. The geometric Newton identity and its exact truncation/remainder are [\[kernel-verified Lean\]](#) in `GeometricNewtonInterpolation`, with endpoints `qNewtonExpansion_eq_pow` and `qNewtonTruncValue_first_omitted`; all broader Bhargava-ordering, Hermite–Padé, and integer-Chebyshev proposals remain research interfaces. No dated literature-search claim is promoted, and none of these statements proves the conjecture.

80 Twentieth report: operator coherence and local-to-global arithmetic

Report 20 is an independent operator-theoretic sibling from the same unified base. It asks which exact operator extensions would add arithmetic information beyond the two scalar values 2^x and 3^x . Bilateral shifts, two-mode thermodynamics, semisimple functional calculus, and bounded Sturmian cocycles are proved blind; global or finite valuation lattices, one affine coherence, one rational first jet, or one exact finite realization would close the problem. These are sufficient criteria and no-go theorems, not constructions supplied by the original hypotheses.

80.1 Universal two-mode dynamics

On $\ell^2(\mathbb{Q}_{>0})$ let $U_a e_q = e_{aq}$. If a, b are multiplicatively independent, the pair (U_a, U_b) is a countable multiple of the regular representation of $\mathbb{Z}^2$: decompose $\mathbb{Q}_{>0}$ into cosets of $\langle a, b \rangle$. For the substantive case $x > 0$, the pairs attached to $(2, 3)$ and $(2^x, 3^x)$ are therefore unitarily equivalent. On the two-mode Fock space,

$$H_{2^x, 3^x} = x H_{2, 3}, \quad (439)$$

so heat traces, local KMS data, and Weyl laws are merely reparametrized in time. This reconciles with report 18: logarithmic spectral rank detects one versus two generators, but cannot quantize a positive rescaling inside the rank-two sector.

The same blindness holds for every rational semisimple matrix T with $(T - 2I)(T - 3I) = 0$. Writing $M = 2^x$ and $A = 3^x$,

$$T^x = (A - M)T + (3M - 2A)I, \quad (440)$$

which is rational, or integral when the displayed data are integral. Coupling the distinct spectral points probes only their already-known secant.

80.2 Global and finite arithmetic covariance

On $\ell^2(\mathbb{N})$ let $H e_n = (\log n) e_n$. For $c > 0$ the following are equivalent:

$$\exists \text{ an isometry } T \text{ with } HT = cTH; \quad n^c \in \mathbb{N} \ (n \in \mathbb{N}); \quad c \in \mathbb{N}. \quad (441)$$

Spectral simplicity gives the middle condition, while the forward difference of order $\lfloor c \rfloor + 1$ is a nonzero integer tending to zero unless c is integral. Conversely $Te_n = e_{nc}$ works for positive integral c . This is much stronger than the original hypotheses, whose canonical partial intertwiner is supported only on the 3-smooth sector.

A finite prime lattice already suffices. Let $P = \{p_1, \dots, p_d\}$ contain 2, assume $p_j^c \in \langle P \rangle$ for every j , and suppose 2^c is algebraic. The valuation columns form $V \in M_d(\mathbb{Z})$ with

$$V^T (\log p_i)_i = c (\log p_i)_i. \quad (442)$$

Thus c is an algebraic integer; Gelfond–Schneider and algebraicity of 2^c make it rational, hence integral. If the lattice map is onto, V is unimodular and positive c must equal 1. Finiteness and exact valuation support are the arithmetic input; an abstract Hilbert-space extension does not provide them.

80.3 Affine rulers and the first-jet criterion

On $\ell^2(\mathbb{N}_0)$ let $Se_k = e_{k+1}$ and $V_a e_k = e_{ak}$. For positive integers a, b, d ,

$$D_{a,b,d} = V_a S^d - S^{bd} V_a \quad (443)$$

vanishes exactly when $a = b$; otherwise every basis vector is sent to the difference of two distinct basis vectors, so $D_{a,b,d}$ is noncompact and $\|D_{a,b,d}\|_{\text{ess}} \geq \sqrt{2}$. Hence preserving the single relation $V_2 S = S^2 V_2$ after $V_2 \mapsto V_{2^x}$, even modulo compacts, forces $2^x = 2$ and $x = 1$. Translation acts as an external integral ruler absent from the multiplicative subsystem.

At one positive repeated spectral point the terminating functional calculus gives, for $b > 0$ and $N^2 = 0$,

$$(bI + N)^x = b^x I + x b^{x-1} N. \quad (444)$$

If b^x and one nonzero first-jet coefficient $x b^{x-1}$ are rational, then $x = b(x b^{x-1})/b^x$ is rational. Under the original integral-output hypotheses this forces $x \in \mathbb{Z}$. Distinct-node Loewner entries are secants and automatic; a self-extension is the first finite-dimensional datum that sees the missing derivative.

80.4 Sturmian similarity and exact finite realization

Writing $x = n + \alpha$, define the rotation jump $w(t) = \lfloor t + \alpha \rfloor$ on $[0, 1)$. The rational two-valued cocycles

$$c_b(t) = b^{w(t)-\alpha} \quad (b = 2, 3) \quad (445)$$

are bounded coboundaries because $w(t) - \alpha = t - \{t + \alpha\}$. Their weighted Koopman operators are therefore similar to the rotation itself; spectra, entropy, Lyapunov growth, and determinant-like invariants are blind. The surviving projection $P = M_w$ has trace $\tau(P) = \alpha$. An exact finite-dimensional trace model would make α rational. Likewise, finite Hankel rank of $w_k = \lfloor (k+1)\alpha \rfloor - \lfloor k\alpha \rfloor$ would make this binary sequence eventually periodic and its mean rational. Either exact realization closes the conjecture; ordinary finite-dimensional approximation does not. The deterministic carry identities overlap the accepted `SturmianProbability` and `SturmianGeneratingSeries` modules.

Status: Report 20 was audited independently from the common base. Its bilateral universality, time-homothety, global and finite covariance criteria, affine essential-norm rigidity, semisimple blindness, Jordan first-jet criterion, and finite trace/Hankel criteria are self-contained paper arguments with the named standard operator inputs. The maximal algebraic support is already covered more generally by `AlgebraicBaseSpectrum`, and the Sturmian/heat cores overlap earlier accepted modules. The rational scalar and matrix-entry first-jet recovery is [\[kernel-verified Lean\]](#) in `JordanFirstJet`; the theorem `integer_of_two_rpow_integer_of_jordanFirstJet_rational` formalizes the arithmetic closing step once the square-zero holomorphic functional-calculus formula is supplied. No global intertwiner, finite prime completion, affine coherence, or finite realization is derived from the two original outputs. None of these criteria proves the conjecture.

81 Varying the antecedent, the domain, and the codomain

The conjecture fixes three choices that look incidental: the values are asked to be *integers*, the exponent is asked to be *real*, and the codomain is $\mathbb{Z}$ rather than a larger ring. This section records what each choice is worth. The short answer is that the conjecture predicts the same thing in every reading; what changes is only which wall stands guard.

81.1 The antecedent ladder

Write the three antecedents

$$(\mathbb{A}_2) \ 2^x, 3^x \in \overline{\mathbb{Q}} \Rightarrow x \in \mathbb{Q}, \quad (\mathbb{Q}_2) \ 2^x, 3^x \in \mathbb{Q} \Rightarrow x \in \mathbb{Z}, \quad (\mathbb{Z}_2) \ 2^x, 3^x \in \mathbb{Z} \Rightarrow x \in \mathbb{Z}.$$

The three-base theorem decomposes into two fully decoupled mechanisms, and only one of them ever sees the difference between $\mathbb{Z}$, $\mathbb{Q}$ and $\overline{\mathbb{Q}}$.

Mechanism one: transcendence. If $2^x, 3^x, 5^x \in \overline{\mathbb{Q}}$ and $x \notin \mathbb{Q}$ then $\{1, x\}$ is $\mathbb{Q}$ -linearly independent, $\{\log 2, \log 3, \log 5\}$ is $\mathbb{Q}$ -linearly independent by unique factorization, and the six products exponentiate to $2, 3, 5, 2^x, 3^x, 5^x$, all algebraic, contradicting six exponentials. This step is blind to integrality: it uses only algebraicity. It is the entire transcendence content of the three-base theorem.

Mechanism two: descent from $\mathbb{Q}$ to $\mathbb{Z}$. If the values are rational, write $x = p/q$ in lowest terms. Then $2^{p/q} \in \mathbb{Q}$ is a rational root of the monic $X^q - 2^p$, hence an integer m , and comparing 2-adic valuations in $2^p = m^q$ gives $p = q v_2(m)$, so $q \mid p$ and $q = 1$. One rational value at one base kills the denominator, and the same works for values in any subring of $\mathbb{Q}$. The dividing line is simply whether the value ring sits inside $\mathbb{Q}$ or contains irrational algebraic numbers.

Status: [\[classical/open background\]](#) for mechanism one; [\[kernel-verified Lean\]](#) for mechanism two, which is `integer_of_rational_of_two_rpow_integer` in `IntegerExponent`. The implications $(\mathbb{A}_2) \Rightarrow (\mathbb{Q}_2) \Rightarrow (\mathbb{Z}_2)$ follow, the first by mechanism two and the second trivially.

The $\overline{\mathbb{Q}}$ -conclusion is sharp. Every rational $x = p/q$ genuinely occurs, since $2^{p/q}, 3^{p/q}, 5^{p/q}$ are algebraic of degree exactly q : the binomial $X^q - 2^p$ is irreducible by Capelli's criterion for $\gcd(p, q) = 1$,

because 2^p is positive and is not an ℓ -th power for any $\ell \mid q$. So the $\overline{\mathbb{Q}}$ -antecedent solution set is exactly $\mathbb{Q}$, and no strengthening of six exponentials can improve the conclusion. The same degree computation gives a graded refinement: if $2^x, 3^x, 5^x$ are algebraic of degree at most d , then $x \in \mathbb{Q}$ with denominator at most d . The kernel-verified exact radical degree of `RadicalDegree` is the two-base conditional shadow of exactly this statement.

Remark 81.1 (Where primality is used). Mechanism one needs only that the bases be multiplicatively independent, but mechanism two needs a base that is not a perfect power. The bases 4, 9, 25 are multiplicatively independent, yet $x = 1/2$ gives $4^x = 2$, $9^x = 3$, $25^x = 5$, all integers. For general multiplicatively independent integer bases the correct statement is: rational values force $x \in \mathbb{Q}$ with a denominator q such that every base is a perfect q -th power. This is exactly the phenomenon that the power-primitive normalization of Section 5 exists to quotient away.

The rational form $(\mathbb{Q}_2)$ deserves separate billing: it is the conjecture as Alaoglu and Erdős posed it, and it is the form the colossally abundant application consumes. Its solution set is exactly $\mathbb{Z}$, negative exponents included, which restores the $x \leftrightarrow -x$ symmetry that the $\mathbb{Z}$ form breaks by truncating to $\mathbb{N}_0$.

Proposition 81.2 (The exact cost of the rational form). *$(\mathbb{Q}_2)$ is equivalent to $(\mathbb{Z}_2)$ together with a denominator-support lemma: if $2^x = a/b$ and $3^x = c/d$ in lowest terms, then b is a power of 2 and d is a power of 3.*

Proof. Given the support lemma, $y = x + s$ for large s has $2^y = 2^s a/b \in \mathbb{Z}$ and $3^y = 3^s c/d \in \mathbb{Z}$, with y nonintegral exactly when x is, so $(\mathbb{Z}_2)$ applies. Conversely $(\mathbb{Q}_2)$ makes the support lemma vacuous, every solution being integral. $\square$

Status: [\[paper proof in this report\]](#) — absent the support lemma no reduction is known in either direction. It is the same kind of statement as the kernel-verified rational gateway of `RationalThirdBase`, and the modules `ThreeDenominatorNormalization` and `RationalPowerIndex` operate inside exactly this gap without stating the equivalence.

81.2 Complexifying the exponent is free

On the principal branch the restriction to real exponents is cosmetic.

Theorem 81.3 (Every complex solution is real). *Let $x \in \mathbb{C}$ and let $2^x = \exp(x \log 2)$, $3^x = \exp(x \log 3)$ be the principal determinations. If $2^x, 3^x \in \mathbb{Z}$ then $x \in \mathbb{R}$. The same holds with $\mathbb{Z}$ replaced by $\mathbb{Q}$.*

Proof. Write $x = s + it$. The modulus of 2^x is 2^s , which reproduces the real hypothesis, and its argument is $t \log 2$; an integer value is real, so $t \log 2 \in \pi\mathbb{Z}$, and likewise $t \log 3 \in \pi\mathbb{Z}$. Write $t \log 2 = \pi a$ and $t \log 3 = \pi b$ with $a, b \in \mathbb{Z}$. If $t \neq 0$ then $a \neq 0$, and dividing gives $\vartheta = b/a \in \mathbb{Q}$, contradicting irrationality of ϑ . $\square$

Status: [\[kernel-verified Lean\]](#) for the integer-codomain version — `im_eq_zero_of_complexTwoBaseIntegralSolution` and `twoBaseIntegralSolution_re_of_complex` in `ComplexSolutions`. The rational-codomain realness statement follows by the same paper argument but is not exported there. Only irrationality of ϑ enters; no transcendence input is used.

Geometrically the sets $\{x : 2^x \in \mathbb{Z}\}$ and $\{x : 3^x \in \mathbb{Z}\}$ lie on horizontal gratings with spacings $\pi/\log 2$ and $\pi/\log 3$, and the two spacings are incommensurable *precisely because* ϑ is irrational, so the gratings meet only on the real axis. The conjecture may therefore be stated over $\mathbb{C}$ at no mathematical cost.

81.3 The algebraic antecedent over the complex numbers

Here the reduction to $\mathbb{R}$ genuinely fails — algebraic values need not have argument in $\{0, \pi\}$ — but something cleaner is true. Since $\overline{\mathbb{Q}}$ is closed under conjugation and division, $z \in \overline{\mathbb{Q}} \setminus \{0\}$ has $|z| = \sqrt{z\bar{z}} \in \mathbb{Q}$. So the complex solution set splits as a direct sum

$$\{x \in \mathbb{C} : 2^x, 3^x \in \overline{\mathbb{Q}}\} = \{s \in \mathbb{R} : 2^s, 3^s \in \overline{\mathbb{Q}}\} + i \{t \in \mathbb{R} : 2^{it}, 3^{it} \in \overline{\mathbb{Q}}\}.$$

The real slice is $(\mathbb{A}_2)$; the imaginary slice is the unit-circle problem. Both are four-exponentials instances — genuinely different instances, with rows $\{1, s\}$ and $\{1, it\}$ against the same columns $\{\log 2, \log 3\}$ — so the complex two-base algebraic problem is exactly two decoupled instances of the same conjecture, and under it the conclusion is still $x \in \mathbb{Q}$, still sharp.

For three bases everything closes unconditionally: six exponentials is a statement over $\mathbb{C}$, the real slice gives $s \in \mathbb{Q}$, and the imaginary slice is the standard corollary that $2^{it}, 3^{it}, 5^{it}$ cannot all be algebraic for real $t \neq 0$. The three-base algebraic solution set over $\mathbb{C}$ is therefore exactly $\mathbb{Q}$; complexification adds nothing. One unconditional crumb survives for two bases: nonreal *algebraic* x are excluded, because Gelfond–Schneider holds for complex algebraic irrational exponents, so the transcendence theorem of Section 4 extends to $\mathbb{C}$ verbatim.

81.4 Multivalued branches widen the wall by one prime

Demand only that *some* determination of each power be an integer: $e^{xL_2} = m$ and $e^{xL_3} = m'$ with $L_2 = \log 2 + 2\pi iJ$ and $L_3 = \log 3 + 2\pi iK$. Writing $xL_2 = \log |m| + i\pi a$ and $xL_3 = \log |m'| + i\pi b$ and eliminating x , the imaginary part gives a $\mathbb{Z}$ -linear relation among logarithms, hence by unique factorization the multiplicative constraint $|m|^{2K} 3^a = |m'|^{2J} 2^b$, while the real part gives

$$\log |m| \log 3 - \log |m'| \log 2 = 2\pi^2 (aK - bJ). \quad (446)$$

If $aK - bJ = 0$ the left side vanishes, which is the log-product wall of Section 12: either $|m| = |m'| = 1$, forcing $x = 0$, or $(|m|, |m'|)$ is a real counterexample pair. If $aK - bJ \neq 0$ one needs the log-product form to equal a nonzero integer multiple of $2\pi^2$. Since $i\pi = \log(-1)$, this is exactly the same wall over the prime set extended by the “prime” -1 : the bilinear lattice gains the directions $\pi^2 = -(\log(-1))^2$ and $i\pi \log p$.

Status: [paper proof in this report] — the independence criterion of Section 12, extended to this larger lattice, covers the multivalued reading as well, and it remains a consequence of Schanuel’s conjecture. The four exponentials conjecture masters the multivalued reading wholesale, because the branch logarithms stay $\mathbb{Q}$ -independent: it gives $x \in \mathbb{Q}$, and then the modulus argument forces $x \in \mathbb{Z}$ with all phases trivial. The predicted answer is unchanged; only the wall widens.

81.5 Complexifying the codomain: the Gaussian version is equivalent

Since extending the domain is free, the honest place to look for new content is the codomain.

Theorem 81.4 (Gaussian equivalence). *The statement “ $2^x, 3^x \in \mathbb{Z}[i] \Rightarrow x \in \mathbb{Z}$ ” is equivalent to Conjecture 2.1.*

Proof. One direction is trivial since $\mathbb{Z} \subset \mathbb{Z}[i]$. Conversely assume the conjecture and let $2^x, 3^x \in \mathbb{Z}[i]$ with $x = s + it$. Taking norms, $N(2^x) = 2^{2s}$ and $N(3^x) = 3^{2s}$ are positive integers, which is the real hypothesis at exponent $2s$, so $2s \in \mathbb{Z}$. Now 3 is inert in $\mathbb{Z}[i]$, so a Gaussian integer of norm 3^{2s} is a unit times 3^k with $9^k = 3^{2s}$, forcing $s = k \in \mathbb{Z}$: the inert prime upgrades the half-integer ambiguity to integrality. Then $2^{2s} = 4^s$ shows that 2^x is divisible only by the ramified prime $1 + i$, so $2^x = u(1 + i)^{2s} = 2^s i^s u$. In both bases the values are 2^s and 3^s times fourth roots of unity, the only Gaussian integers of modulus one, so the argument conditions read $t \log 2, t \log 3 \in (\pi/2)\mathbb{Z}$, and irrationality of ϑ kills t as in Theorem 81.3. $\square$

Status: [\[paper proof in this report\]](#) — a self-contained equivalence whose only inputs are Gaussian factorization, irrationality of ϑ , and the real conjecture as a hypothesis. It is a low-cost Lean target on top of `ComplexSolutions`.

Two boundary markers make the picture sharp. Replacing $\mathbb{Z}[i]$ by $\mathbb{Q}(i)$ breaks the argument at exactly one point: unit-modulus elements of $\mathbb{Q}(i)$ need not be roots of unity, as $(3 + 4i)/5$ shows, so the imaginary slice reopens as a genuine unit-circle four-exponentials instance. It is integrality up to units, not rationality, that closes the argument direction. Replacing $\mathbb{Z}[i]$ by the integers of an imaginary quadratic field in which 2 or 3 *splits* — for instance $\mathbb{Q}(\sqrt{-2})$, where 3 splits and $1 + \sqrt{-2}$ has modulus $\sqrt{3}$ — destroys the root-of-unity conclusion, and eliminating t produces relations mixing $\vartheta\pi$ with logarithms of complex algebraic numbers: the product-of-logarithms wall again, with more rooms. These split-prime variants are recorded as open rather than pursued; they meet the same wall with worse constants.

81.6 Summary

Reading	Status
Exponent in $\mathbb{C}$, principal branch, values in $\mathbb{Z}$	reduces to the real integer formulation, [kernel-verified Lean]
Exponent in $\mathbb{C}$, principal branch, values in $\mathbb{Q}$	reduces to the real rational-output formulation at paper level; equivalence to integer AE still needs the denominator-support lemma of Proposition 81.2
Values in $\overline{\mathbb{Q}}$, exponent real	two-base: four exponentials; conclusion $\mathbb{Q}$, sharp
Values in $\overline{\mathbb{Q}}$, exponent in $\mathbb{C}$	splits into real and unit-circle slices
Three bases, values in $\overline{\mathbb{Q}}$, exponent in $\mathbb{C}$	solution set exactly $\mathbb{Q}$, [classical/open background]
Multivalued branches	wall widened by $\log(-1)$, [paper proof in this report]
Values in $\mathbb{Z}[i]$	unconditionally equivalent, [paper proof in this report]
Values in $\mathbb{Q}(i)$ or split-prime rings	open

In every reading the answer the conjecture predicts is the same.

82 Twenty-first report: Hausdorff moments and Müntz invisibility

Report 21 is independent of the previous continuations. Its minimal transcendence and fractional-normalization steps overlap the accepted corpus; its new direction is the additive adjacency of the $\{2, 3\}$ -smooth numbers and the Hausdorff moment problem it creates.

82.1 The explicit moment measure

For $0 < \alpha < 1$ define

$$d_j = (j+2)^\alpha - (j+1)^\alpha \quad (j \geq 0).$$

These are the moments of the finite positive measure

$$d\mu_\alpha(q) = \frac{\alpha}{\Gamma(1-\alpha)}(1-q)(-\log q)^{-\alpha-1} dq, \quad d_j = \int_0^1 q^j d\mu_\alpha(q). \quad (447)$$

The same Mellin identity holds for every real exponent $z > -1$. It follows from the gamma integral for $t^{\alpha-1}$ and Tonelli after $q = e^{-s}$. The density is positive on $(0, 1)$, integrable at both endpoints, and has full support. Consequently

$$(-1)^k \Delta^k d_j = \int_0^1 q^j (1-q)^k d\mu_\alpha(q) > 0, \quad (448)$$

every finite Hankel matrix (d_{i+j}) is positive definite, and strict Hölder gives

$$d_j^{k-i} < d_i^{k-j} d_k^{j-i} \quad (0 \leq i < j < k).$$

These analytic identities are [\[paper proof in this report\]](#) paper proofs using classical integration.

The only adjacent positive $\{2, 3\}$ -smooth pairs are

$$(1, 2), \quad (2, 3), \quad (3, 4), \quad (8, 9),$$

by coprimality and Catalan–Mihăilescu. Thus under a hypothetical fractional solution, with $u = 2^\alpha$ and $v = 3^\alpha$, the adjacent endpoints force

$$d_0 = u - 1, \quad d_1 = v - u, \quad d_2 = u^2 - v, \quad d_7 = v^2 - u^3 \quad (449)$$

to be rational and positive. This is a list of the four moments whose rationality is *forced by adjacent smooth endpoints*; it does not exclude accidental rationality of other moments.

82.2 Exact arithmetic defects and finite rational simulation

Write $u = P/2^r$ and $v = Q/3^s$ in lowest terms. The first Hankel defect and three nonconsecutive log-convexity defects are

$$\mathcal{H} = d_0 d_2 - d_1^2, \quad \mathcal{F}_{017} = d_0^6 d_7 - d_1^7, \quad \mathcal{F}_{127} = d_1^5 d_7 - d_2^6, \quad \mathcal{F}_{027} = d_0^5 d_7^2 - d_2^7.$$

They are strictly positive, and their exact reduced denominators are respectively

$$2^{3r} 3^{2s}, \quad 2^{9r} 3^{7s}, \quad 2^{12r} 3^{7s}, \quad 2^{14r} 3^{7s}. \quad (450)$$

Parity and reduction modulo 3 show that the displayed positive numerators are coprime to these denominators. The four common-denominator identities and their local-unit conclusions are [\[kernel-verified Lean\]](#) in `Report21MomentArithmetic`; positivity comes from (447) and remains the paper layer.

Finite positivity data cannot distinguish the special measure. If rational $m_0 > m_1 > m_2 > 0$ satisfy $m_0 m_2 > m_1^2$, then

$$\left(m_0 - \frac{m_1^2}{m_2}\right) \delta_0 + \frac{m_1^2}{m_2} \delta_{m_2/m_1} \quad (451)$$

is a positive rational two-atom measure reproducing m_0, m_1, m_2 . This identity is also [\[kernel-verified Lean\]](#) in `Report21MomentArithmetic`. More generally, a finite family of rational polynomial observations lying in the interior of the closed moment cone admits a positive rational finite quadrature: one first passes to the closed cone, uses $\text{int } C = \text{int } \overline{C}$, and then perturbs a finite generating set to rational nodes. This is a [\[paper proof in this report\]](#) finite-dimensional paper theorem.

82.3 All smooth linear observations remain invisible

Put

$$\Lambda = \{N - 1 : N = 2^a 3^b \geq 2\}, \quad K_m(q) = \frac{1 - q^m}{1 - q}.$$

Unique factorization gives the convergent reciprocal sum

$$\sum_{m \in \Lambda} \frac{1}{m} \leq 4. \quad (452)$$

The full L^1 Müntz theorem on an interval away from zero, followed by Hahn–Banach, now gives an exact infinite-dimensional no-go. For every $0 < a < b < 1$ there is a nonzero bounded real g , supported in $[a, b]$, such that

$$\int_0^1 K_m(q) g(q) d\mu_\alpha(q) = 0 \quad (m \in \Lambda). \quad (453)$$

For sufficiently small $\varepsilon > 0$, the distinct positive measures $(1 \pm \varepsilon g)\mu_\alpha$ therefore have every same $\{2, 3\}$ -smooth linear observation, agree with μ_α near both endpoints, and change the j th ordinary moment by only $O(b^j)$. This is a [\[paper proof in this report\]](#) theorem conditional only on the classical Müntz/Hahn–Banach inputs. It is a no-go for reconstruction among positive measures; the perturbations need not belong to the one-parameter family $\{\mu_\beta\}$ and therefore do not produce another exponent.

Status: Report 21 was audited independently. The exact two-atom and denominator arithmetic is [\[kernel-verified Lean\]](#) in `Report21MomentArithmetic`. The Hausdorff integral, finite moment-cone quadrature, and Müntz invisible perturbations are paper/classical arguments with the scopes stated above. They prove that positivity and all smooth linear observations are insufficient, not that the conjecture fails or holds.

83 Twenty-second report: depth-two rank geometry and rational symmetrization

Report 22 is another independent sibling of the common base. Its restricted matrix coefficient, formal prime-symbol, and external-rank reductions overlap `RestrictedMatrixCoefficient`, `FormalPrimeDeterminant`, and the accepted external-rank modules. The genuinely new layer is a quadratic-rank classification and a rational symmetric normal form for the multiplicatively dependent branch.

83.1 Baker excludes quadratic rank at most two

Let $\lambda_1, \dots, \lambda_t$ be $\mathbb{Q}$ -independent logarithms of positive algebraic numbers. A nonzero rational quadratic form q of rank at most two cannot vanish at these logarithms. Over $\overline{\mathbb{Q}}$, such a form factors as $q = L_1 L_2$; vanishing would give a nontrivial algebraic linear relation among the λ_i , contrary to Baker’s homogeneous linear-independence theorem. This is a classical Baker consequence; its application to the two-star tensor is new to this corpus.

In prime-exponent space put $a = \nu(M)$, $b = \nu(N)$ and

$$q_{a,b} = e_2 b - e_3 a \in \text{Sym}^2(V).$$

Its period specialization is $\log 2 \log N - \log 3 \log M$. A nonintegral solution makes this zero but does not make the formal tensor zero, so Baker forces $\text{rank}(q_{a,b}) \geq 3$.

83.2 Hyperbolic pullback and the rank-three conic

For independent e, f and arbitrary a, b , the universal tensor $eb - fa$ is the pullback of the hyperbolic form $z_1 z_2 - z_3 z_4$. If $r = \dim \text{span}\{e, b, f, a\}$, its rank is four for $r = 4$ and at most two for $r \leq 2$. At $r = 3$, write the unique relation as

$$c_1 e + c_2 b + c_3 f + c_4 a = 0.$$

The restricted form has rank three exactly when

$$c_1 c_2 - c_3 c_4 \neq 0; \tag{454}$$

otherwise its rank is two. Combining this finite classification with Baker shows that a hypothetical counterexample has equal multiplicative and quadratic rank, necessarily three or four. In rank three the two external prime-exponent vectors lie on one rational ray, and the determinant relation becomes a nondegenerate split ternary conic. This is a compression of the dependent branch, not a nonvanishing theorem for its prime-log point.

The hyperbolic normal identity and exact conic determinant are [\[kernel-verified Lean\]](#) in `Report22RankGeometry`. Baker’s theorem and the passage from the actual two-star tensor to the rank classification are the paper/classical layer. The source report wrote the rank-three and rank-four kernel strata as intersections equal to $\{0\}$; literally those intersections must be empty, because the zero tensor is not in either stratum. The unified formulation uses the corrected nonvanishing statement.

83.3 A rational symmetric normal form

Let $L \in M_2(\mathbb{R})$ be nonzero and singular, suppose its entries span a three-dimensional rational vector space, and let its unique rational coefficient relation have matrix

$$C = \begin{pmatrix} \gamma & \alpha \\ \delta & \beta \end{pmatrix}, \quad \Delta = \gamma\beta - \delta\alpha \neq 0.$$

The explicit right multiplier

$$Q = \begin{pmatrix} -\delta & \gamma \\ -\beta & \alpha \end{pmatrix}, \quad \det Q = \Delta, \quad (455)$$

makes LQ symmetric. If $L = r(1, x)$ has rank one, symmetry forces

$$LQ = \kappa r r^T.$$

For a rank-three counterexample, $r = (\log 2, \log 3)^T$; after changing the global sign, $\kappa > 0$, and the three entries

$$\kappa(\log 2)^2, \quad \kappa \log 2 \log 3, \quad \kappa(\log 3)^2$$

are independent logarithms of positive algebraic numbers in geometric progression. Thus there would be a positive algebraic $\eta > 1$ for which

$$\eta, \quad \eta^\vartheta, \quad \eta^{\vartheta^2} \quad \text{are algebraic and multiplicatively independent.} \quad (456)$$

Current transcendence theorems do not exclude this two-step orbit. The determinant, right-multiplication symmetry, invertible-multiplier construction, and symmetric outer-product normal form are [\[kernel-verified Lean\]](#) in `Report22RankGeometry`; algebraicity and the application to the counterexample remain the [\[paper proof in this report\]](#) paper layer.

Status: Report 22 was audited independently. Its low-rank exclusion is a classical Baker consequence, and its hyperbolic pullback and rational symmetrization are new exact reductions. The finite matrix identities are kernel-checked in `Report22RankGeometry`. The remaining rank-three and rank-four depth-two period nonvanishing statements are open; neither the conic nor the symmetric orbit proves the conjecture.

84 Twenty-third report: unimodular gaps, transported secants, and finite-order barriers

Report 23 is an independent continuation from the common unified base. Its elementary and Four/Six Exponentials reductions coincide with accepted material and are not repeated. Its new contribution is a determinant-distance law for binomial gaps, an associated family of rational secant brackets, and two precise no-go theorems for finite order data and bounded entrywise jets.

84.1 Determinant distance and adjacent gap coprimality

Let $A, B \geq 2$ be coprime, let p, q, r, s be positive integers, and put $D = |ps - rq| > 0$. Then

$$\gcd(B^q - A^p, B^s - A^r) \mid \gcd(A^D - 1, B^D - 1). \quad (457)$$

Indeed a common divisor is prime to AB ; eliminating first B and then A from the two power congruences gives $A^D \equiv B^D \equiv 1$. When $D = 1$ the converse congruence is also automatic, so the exact identity is

$$\gcd(B^q - A^p, B^s - A^r) = \gcd(A - 1, B - 1). \quad (458)$$

Without coprimality, the same statement holds primewise away from AB . Consecutive continued-fraction convergents to $\log 3 / \log 2$ have determinant distance one; hence their signed gaps $3^{q_j} - 2^{p_j}$ are pairwise coprime at adjacent indices. For a hypothetical output pair $(M, N) = (2^x, 3^x)$, the corresponding adjacent output-gap gcd is $\gcd(M - 1, N - 1)$ when M, N are coprime, with the same away-from- MN restriction otherwise. These are [\[paper proof in this report\]](#) exact paper theorems. The generic commutative-group collision and both determinant orientations of the modular implication are [\[kernel-verified Lean\]](#) in `UnimodularGapCollision`; the integer gcd packaging remains the paper layer.

84.2 Rational secants and separated denominators

For positive p, q with $3^q \neq 2^p$, define

$$Q_{p,q} = \frac{N^q - M^p}{3^q - 2^p}.$$

Under a hypothetical solution $x > 1$ this is positive rational and the fundamental theorem of calculus gives

$$Q_{p,q} = x \int_0^1 ((1-t)2^p + t3^q)^{x-1} dt. \quad (459)$$

Its reduced denominator is exactly $|3^q - 2^p| / \gcd(|3^q - 2^p|, |N^q - M^p|)$. Thus adjacent convergents give coprime reduced denominators. After setting

$$T_j = \frac{2^{p_j}}{M^{p_j}} Q_{p_j, q_j}, \quad \delta_j = q_j \log 3 - p_j \log 2,$$

one has

$$T_j = \Phi_x(\delta_j), \quad \Phi_x(\delta) = \frac{e^{x\delta} - 1}{e^\delta - 1}, \quad T_j - x = \frac{x(x-1)}{2} \delta_j + O_x(\delta_j^2).$$

The integral representation of Φ_x makes it strictly increasing, so the rational T_j alternate above and below x . Away from primes dividing M , consecutive reduced denominators remain coprime. This is exact arithmetic separation, not an irrationality measure: the error is only $O(1/q_{j+1})$ while the rational heights can be exponential in p_j .

84.3 Finite signatures and the entrywise jet ceiling

No finite list of strict monomial comparisons identifies the special slope. For every finite $F \subset \mathbb{Z}^2 \setminus \{0\}$ there are infinitely many coprime multiplicatively independent pairs (M, N) such that

$$\operatorname{sgn}(a \log M + b \log N) = \operatorname{sgn}(a \log 2 + b \log 3) \quad ((a, b) \in F).$$

Choose a large prime $M = P$ and an integer N nearest $P^{\log_2 3}$ but not divisible by P ; the log slope lies in the same finite sign chamber. This rules out only finite strict order signatures, not congruence, valuation, or uniform infinite information.

There is a complementary determinant ceiling. For an $R \times R$ integer decoration matrix C , let $\Sigma(C)$ be the set of permutations whose Leibniz products are nonzero and assume $\Sigma(C) \neq \emptyset$. Let $\tau_p(C)$ be the minimum p -valuation over $\Sigma(C)$ and put $T(C) = \prod_p p^{\tau_p(C)}$. If every nonzero $|c_{ij}| \leq B$, then for every integral E ,

$$T(C) \mid \det(E \circ C), \quad \log T(C) \leq R \log B. \quad (460)$$

For Euler or Hasse multipliers of order at most J on exponent scale L , this standalone guaranteed divisor contributes at most $RJ \log L$. It need not be coprime to, or additive with, existing determinant divisors. Therefore only the entrywise/termwise multiplier contribution is $o(R^2)$ when $J \log L = o(R)$; global confluence, cancellation, or factorization remains outside the theorem. The gap, secant, signature, and decoration results are [\[paper proof in this report\]](#) paper arguments and do not prove the conjecture.

Status: Report 23 was reconciled as an independent sibling. The genuinely new determinant-distance collision core is [\[kernel-verified Lean\]](#) in `UnimodularGapCollision`; the continued-fraction gcd, transported-secant, finite-signature, and entrywise jet packages are paper proofs. Existing transcendence reductions are cross-referenced, and no fixed-order jet statement is promoted to a claim about global determinant cancellation.

85 Twenty-fourth report: exact orbit heights and Farey entropy

Report 24 is another independent continuation from the common unified base. Its compact orbit, recurrence, contact-rigidity, and Four Exponentials ingredients overlap earlier sections. Its genuinely new contribution is quantitative: it computes every reduced orbit height exactly, converts Farey counting into a simultaneous two-coordinate denominator lower bound, and compares Baker recurrence with a fixed analytic contact order.

85.1 Exact reduced orbit law

Condition on a nonintegral counterexample and write $x = B + \alpha$ with $0 < \alpha < 1$. In lowest terms put

$$2^\alpha = \frac{W}{2^r}, \quad 3^\alpha = \frac{Z}{3^s}, \quad W > 1 \text{ odd}, \quad 3 \nmid Z.$$

Thus $r + \alpha = \log_2 W$ and $s + \alpha = \log_3 Z$. For $k \geq 0$ let $(R_k, S_k) = (2^{\{k\alpha\}}, 3^{\{k\alpha\}})$, and for a reduced rational a/b set $H(a/b) = \max(|a|, b)$. For every $k \geq 1$ the exact reduced expressions are

$$R_k = \frac{W^k}{2^{\lfloor k \log_2 W \rfloor}}, \quad S_k = \frac{Z^k}{3^{\lfloor k \log_3 Z \rfloor}}, \quad (461)$$

so

$$H(R_k) = W^k, \quad H(S_k) = Z^k, \quad \frac{W^k}{2} < \text{den}(R_k) < W^k, \quad \frac{Z^k}{3} < \text{den}(S_k) < Z^k. \quad (462)$$

The floor identity behind this is $rk + \lfloor k\alpha \rfloor = \lfloor k(r + \alpha) \rfloor$; coprimality makes the displayed fractions reduced. Consequently, with $C_0 = \max(W, Z)$,

$$\#\{k \geq 0 : \max(H(R_k), H(S_k)) \leq H\} = 1 + \left\lfloor \frac{\log H}{\log C_0} \right\rfloor \quad (H \geq 1). \quad (463)$$

The floor and finite simultaneous-denominator cores are [\[kernel-verified Lean\]](#) in `OrbitHeightEntropy`; the reduced-fraction and height specialization is the [\[paper proof in this report\]](#) paper layer.

85.2 Worst-case Farey-to-orbit distortion

Fix a nondegenerate closed interval $I \subset (0, 1)$ and let $\mathcal{F}_I(Q)$ be its reduced rationals with denominator at most Q . The classical estimate

$$\#\mathcal{F}_I(Q) = \frac{3|I|}{\pi^2} Q^2 + O_I(Q \log Q)$$

combines with (462) as follows. If

$$\Phi : \mathbb{Q} \cap I \longrightarrow \{(R_k, S_k) : k \geq 0\}$$

is injective and $D_2(Q), D_3(Q)$ are the maximum first- and second-coordinate denominators over $\mathcal{F}_I(Q)$, then

$$\liminf_{Q \rightarrow \infty} \frac{\log D_2(Q)}{Q^2} \geq \frac{3|I|}{\pi^2} \log W, \quad \liminf_{Q \rightarrow \infty} \frac{\log D_3(Q)}{Q^2} \geq \frac{3|I|}{\pi^2} \log Z. \quad (464)$$

Indeed N distinct nonnegative orbit indices include one index at least $N - 1$; the *same* image therefore has denominators exceeding $W^{N-1}/2$ and $Z^{N-1}/3$. The argument also permits fibers of a fixed integer multiplicity $M \geq 1$, dividing the two constants by M . The finite injective engine is [\[kernel-verified Lean\]](#) in `OrbitHeightEntropy`; the Farey asymptotic and limiting constants are [\[paper proof in this report\]](#) paper mathematics.

This is a maximum-envelope theorem: for each large cutoff it forces at least one badly distorted input, not every input. It excludes polynomially tame injective, or uniformly bounded-multiplicity, relabelings of the existing orbit. It does not exclude an arithmetic operation that creates new rational graph points. A denominator-tame thickening with polynomially many genuinely new points is therefore a research target, not a consequence of the entropy theorem.

85.3 Baker recurrence versus fixed contact

Let p_j/q_j run through the convergents below α , so $0 < \delta_j = q_j \alpha - p_j < q_j^{-1}$. A two-logarithm Baker estimate [3] gives constants $c, \kappa > 0$ with $\delta_j \geq c q_j^{-\kappa}$. If $0 \neq P \in \mathbb{Q}[X, Y]$, $P(1, 1) = 0$, and

$$\sigma(P) = \text{ord}_{t=0} P(2^t, 3^t),$$

then eventual nonvanishing and Taylor comparison give

$$c_P q_j^{-\kappa \sigma(P)} \leq |P(R_{q_j}, S_{q_j})| \leq C_P q_j^{-\sigma(P)}. \quad (465)$$

The separate elementary height estimate $|F(R_k, S_k)| \geq e^{-C_F k}$ applies when $F \in \mathbb{Q}(X, Y)$ is regular near the compact arc, $k \geq 1$, and $F(R_k, S_k) \neq 0$; this corrects the source statement, which fails at $k = 0$ for constants of absolute value below one. The contact-multiplicity core is already stronger in `RationalContactRigidity`; the new paper contribution here is its two-sided composition with the Baker return scale.

More abstractly, if nonzero rational certificates satisfy

$$\text{den}(\Xi_N) \leq e^{A_N N}, \quad |\Xi_N| \leq C_N N^{-d_N},$$

with $A_N, d_N \geq 0$ and $C_N > 0$, the elementary denominator lower bound can contradict the analytic estimate only if

$$d_N \log N > A_N N + \log C_N. \quad (466)$$

Here “subexponential” means $\log C_N = o(N)$. This is a necessary scale benchmark, not a construction and not a proof of the conjecture.

Status: Report 24 was audited as an independent sibling. The exact floor identities and finite worst-case entropy mechanism are [\[kernel-verified Lean\]](#) in `OrbitHeightEntropy`. The exact rational heights, Farey limits, and Baker–contact comparison are [\[paper proof in this report\]](#) paper results. The entropy theorem concerns relabelings of the existing cyclic orbit; a genuine arithmetic thickening and every closing jet construction remain open.

86 Twenty-fifth report: projective Padé defect and rational-transfer no-go

Report 25 independently tests the new weight-two polylogarithmic criterion of Kawashima [43] against the logarithm determinant. Its tensor-rank reduction overlaps the stronger rank geometry in Section 83, and its directed-rounding scan repeats an unbundled computation already separated from the certified finite frontier. Neither is counted again. The new contribution is an optimized, representation-independent height defect and the no-go theorem it forces for rational polylogarithm points.

86.1 The projective defect bound

For $m \geq 1$ put $D_m = (m + 1)^2 - 1$ and

$$C_m = D_m \log 2 + 3 \log(m + 1) + 2 + 2D_m.$$

Let K be a number field, v a place, let $\alpha_1, \dots, \alpha_m \in K^\times$ be pairwise distinct, and suppose $\beta \in K$ satisfies $|\beta|_v > H_v(\boldsymbol{\alpha}) := \max(1, \max_i |\alpha_i|_v)$. Set $z_i = \alpha_i/\beta$ and write

$$a_v(\mathbf{z}) = -\log \max_i |z_i|_v, \quad \bar{h}(\mathbf{z}) = \frac{1}{m} \sum_{i=1}^m h(z_i).$$

Kawashima’s weight-two defect $\mathcal{V}_v(\boldsymbol{\alpha}, \beta)$ obeys the exact projective upper bound

$$\mathcal{V}_v(\boldsymbol{\alpha}, \beta) \leq (D_m + 1)a_v(\mathbf{z}) - D_m \bar{h}(\mathbf{z}) - C_m. \quad (467)$$

The proof writes $t = h_v(\beta)$ and optimizes

$$(D_m + 1)t - (D_m + 1) \max(0, t - a_v) - D_m \max(t, \bar{h}) - C_m.$$

Since $a_v \leq \bar{h}$, the maximum is attained throughout $a_v \leq t \leq \bar{h}$ and equals the right side of (467). This piecewise optimization and its abstract height-budget wrapper are [\[kernel-verified Lean\]](#) in `Report25ProjectiveDefect`; the number-field height specialization and Kawashima’s theorem are respectively [\[paper proof in this report\]](#) and [\[classical/open background\]](#).

If the defect is positive, then necessarily

$$a_v > C_m, \quad \frac{a_v}{\bar{h}} > \frac{D_m}{D_m + 1} = 1 - \frac{1}{(m+1)^2}, \quad \bar{h} < \left(1 + \frac{1}{D_m}\right) a_v. \quad (468)$$

These implications are also [\[kernel-verified Lean\]](#) in `Report25ProjectiveDefect`. Thus mere archimedean proximity to one is insufficient: the local depth must consume almost all of the global height. Direct points $1 - 1/q$, high-root splittings, and ordinary unit-fraction telescoping all fail this necessary condition. A common integral rescaling is strictly harmful, decreasing the affine defect by $2D_m \log t$.

86.2 No rational transfer of the prime direction

Let distinct $z_1, \dots, z_m \in \mathbb{Q} \cap (0, 1)$ admit a positive-defect specialization at the real place, and put $\ell_i = -\log(1 - z_i)$. Then

$$\log 2 \notin \text{span}_{\mathbb{Q}}\{\ell_1, \dots, \ell_m\}. \quad (469)$$

The same proof applies to a fixed $q \in \mathbb{Q}_{>0} \setminus \{1\}$ with a changed constant. Indeed (468) makes every $\ell_i = O(e^{-a_v})$ while bounding all prime valuations of $(1 - z_i)^{-1}$ by $O(ma_v)$. If the fixed prime direction lay in their rational span, Cramer’s rule and Hadamard’s inequality would supply integer coefficients of size polynomial in a_v ; exponentiation would then force the nonzero constant $\log 2$ below an exponentially decaying bound. The final elementary comparison is uniform in m (with the logarithm taken of the ratio of that bound to $\log 2$ and with the endpoint derivative checked explicitly). This is a [\[paper proof in this report\]](#) paper theorem.

Consequently a positive-defect rational tuple cannot reproduce the counterexample’s formal four-log tensor. The range of that tensor contains the prime directions by Section 83, whereas every tensor built from these rational logarithms has range inside their forbidden valuation span. This rules out the factor-by-factor rational Li_1 transfer; it does not rule out algebraic points, varying fields, or a genuinely mixed-denominator determinant construction.

For a fixed number field K , place v , finitely generated $\Gamma < K^\times$, and fixed cardinality m , the effective approximation theorem of Calegari–Dimitrov–Tang [44] further implies that only finitely many tuples $z_i = 1 - \gamma_i$, $\gamma_i \in \Gamma$, can have positive defect, with effective height bounds. The fixed- m qualification is essential; this is not uniform over constructions with $m \rightarrow \infty$.

Finally, let p_n/q_n be the continued-fraction convergents to $\log 3/\log 2$ and choose $\varepsilon_n \in \{\pm 1\}$ so that the units $r_n = (2^{p_n} 3^{-q_n})^{\varepsilon_n} > 1$; then $r_n \rightarrow 1$ and

$$\frac{-\log(1 - r_n^{-1})}{h(1 - r_n^{-1})} \rightarrow 0,$$

and the output units have the same low-efficiency behavior. The present Padé criterion requires efficiency tending to one. This identifies an arithmetic-efficiency gap, not a closing theorem. A proposed cross-row, mixed-denominator cancellation remains only a schematic research target because its exceptional “zero coefficient after rational changes” alternative has not yet been defined precisely.

Status: Report 25 was audited independently. The finite projective optimization and efficiency implications are [\[kernel-verified Lean\]](#) in `Report25ProjectiveDefect`; the number-field specialization, rational-transfer no-go, and fixed-cardinality group finiteness are [\[paper proof in this report\]](#) paper results using the cited external theorems. The repeated 2^{29} scan is not promoted, and no mixed-denominator construction or proof of the conjecture is claimed.

87 Twenty-sixth report: carry languages and fixed-target entropy

Report 26 is an independent sibling. Its logarithmic cone, one-orbit equivalence, one/two/three-base threshold, 3-smooth lattice count, and residual compact orbit largely overlap the accepted corpus. Its distinct contribution is additive: a hypothetical counterexample forces a distinguished-descendant carry language whose words give many equal-arity representations of one fixed target. Character balancing explains why these families lie in algebraic parts rather than contradicting functional transcendence.

87.1 Carry paths and finite alphabets

Let $U, V \geq 2$ be multiplicatively independent integers and write $e_{a,b} = U^a V^b$. A step decreasing the first exponent replaces one distinguished copy by U copies, freezes $U - 1$, and continues with the remaining copy; the second-coordinate step does the same with V . A monotone word with r first steps and s second steps therefore yields a multiset representation

$$U^r V^s = \sum_{j=1}^{K_{r,s}} e_{a_j, b_j}, \quad K_{r,s} = 1 + r(U - 1) + s(V - 1). \quad (470)$$

The frozen terms telescope to the initial monomial. At every intermediate total degree the multiset records the unique visited vertex; multiplicative independence then reconstructs the word. Hence there are at least

$$\binom{r+s}{r} \quad (471)$$

distinct representations of the fixed target. Under a hypothetical counterexample this applies to $(U, V) = (M, N) = (2^x, 3^x)$ and every term is itself $(2^a 3^b)^x$.

The construction extends to a finite alphabet of distinct moves $v_i = (\alpha_i, \beta_i) \in \mathbb{N}_0^2 \setminus \{(0, 0)\}$. Put

$$D_i = U^{\alpha_i} V^{\beta_i}, \quad (r, s) = \sum_i n_i v_i, \quad K = 1 + \sum_i n_i (D_i - 1).$$

Words with the same symbol counts then give at least

$$\binom{\sum_i n_i}{n_1, \dots, n_m} \quad (472)$$

distinct K -term representations of $U^r V^s$. The exact telescoping sum, term count, strict suffix-degree descent, and reconstruction from the endpoint and ordered suffix list are [\[kernel-verified Lean\]](#) in

Report26CarryPath. Recovering the word from the unordered frozen multiset and the multinomial enumeration remain the [\[paper proof in this report\]](#) paper layer.

The two primitive counterexample carries $(2t)^x = Mt^x$ and $(3t)^x = Nt^x$ are also exactly character-balanced. More generally every finite sequence of such rewrites groups into independently vanishing character classes. Thus their positive-dimensional families are compatible with, rather than forbidden by, the usual functional-transcendence decomposition.

87.2 Exact carry-language pressure

For a fixed finite alphabet $\mathcal{S}$ of at least two 3-smooth integers $d_i = 2^{\alpha_i} 3^{\beta_i} > 1$, let $D_i = d_i^x$ under the hypothetical solution. Entropy maximization gives two exact *carry-language* rates. They are the unique positive roots

$$\sum_i d_i^{-\delta_{\mathcal{S}}} = 1, \quad \sum_i e^{-\tau_{\mathcal{S}}(D_i-1)} = 1. \quad (473)$$

For symbol proportions π_i , relative entropy proves

$$\sup_{\pi} \frac{H(\pi)}{\sum_i \pi_i \log d_i} = \delta_{\mathcal{S}}, \quad \sup_{\pi} \frac{H(\pi)}{\sum_i \pi_i (D_i - 1)} = \tau_{\mathcal{S}}.$$

The alphabet $\{2, 3, 6\}$ already gives ambient exponent one. Exhausting all nonunit 3-smooth letters gives

$$(1 - 2^{-\delta_*})(1 - 3^{-\delta_*}) = \frac{1}{2}, \quad \delta_* = 1.4352790836694176 \dots \quad (474)$$

This is an exact supremum over the finite-alphabet carry construction, not over all additive representations.

There is also an unrestricted word pressure. Let $\tau_x > 0$ be the unique solution of

$$\sum_{(a,b) \in \mathbb{N}_0^2 \setminus \{(0,0)\}} e^{-\tau_x(M^a N^b - 1)} = 1, \quad g = \gcd(M - 1, N - 1). \quad (475)$$

For every sufficiently large positive multiple n of g , some endpoint (r_n, s_n) with $r_n + s_n \leq n$ receives

$$\gg_x \frac{e^{\tau_x n}}{n^2} \quad (476)$$

distinct carry-language representations of the fixed target $M^{r_n} N^{s_n}$ as a sum of $n + 1$ values a^x , all with $a \leq 3^n$. The word generating series is

$$\frac{1}{1 - \Phi_x(z)}, \quad \Phi_x(z) = \sum_{(a,b) \neq (0,0)} z^{M^a N^b - 1}.$$

After factoring the period g , its unique dominant positive pole gives $[z^n](1 - \Phi_x)^{-1} \asymp_x e^{\tau_x n}$ along $g \mid n$; pigeonholing among $O(n^2)$ endpoints yields (476). The matching upper rate is only for fibers generated by this unrestricted carry language, not for the full additive representation fiber.

These are [\[paper proof in this report\]](#) exact paper theorems. They do not conflict with Harrison’s fixed-arity and fixed-target bounds [45]: here the arity grows linearly with the logarithm of the input cutoff. A closing theorem would need uniform fixed-target bounds through that growing-arity range, strong enough to beat $e^{\tau_x n}$; no such estimate is known. Compact-orbit and point-counting reformulations do not supply it.

Status: Report 26 was reconciled independently. The bounded carry-word sum, length, suffix-degree, and ordered reconstruction identities are [\[kernel-verified Lean\]](#) in `Report26CarryPath`. Multiset injectivity, multinomial fibers, character balancing, and the finite and unrestricted pressure laws are [\[paper proof in this report\]](#) paper results. Every “exact rate” is explicitly scoped to the carry-language construction; no contradiction with fixed-arity theorems or proof of Alaoglu–Erdős is claimed.

88 Twenty-seventh report: canonical-product threshold and critical defect cocycle

Report 27 independently studies the power function on the universal cover of $\mathbb{C}^\times$ and the arithmetic sampling set $\{2^a 3^b : a, b \geq 0\}$. Its Six Exponentials monodromy and algebraic-multiplier consequences overlap Sections 65.6 and 76.1; they are not repeated. Its new core is an exact analytic growth threshold for the 3-smooth lattice and an infinite boundary-defect identity at that threshold.

88.1 Regularity ladder for a power branch

On the universal cover, a holomorphic function satisfying

$$F(w + \log 2) = 2^x F(w), \quad F(w + \log 3) = 3^x F(w), \quad F(0) = 1$$

is necessarily $F(w) = e^{xw}$: after division by e^{xw} , the two incommensurable real periods have a dense subgroup, so continuity makes the quotient constant. The corresponding branch z^x descends to a single-valued holomorphic function on $\mathbb{C}^\times$ exactly when $e^{2\pi i x} = 1$, hence exactly when $x \in \mathbb{Z}$. Throughout the algebraic, Puiseux, and D-finite discussion, z^x means a branch on a simply connected punctured neighborhood or on this universal cover, not an ordinary analytic germ at zero.

Finite ramification makes the monodromy a root of unity and forces $x \in \mathbb{Q}$, after which the elementary rational-exponent argument closes the conjecture. A further exact criterion is arithmetic D-finiteness. If a branch z^λ satisfies a nonzero linear differential equation with coefficients in $\overline{\mathbb{Q}}[z]$, of order r and leading coefficient $p_r \neq 0$, then clearing powers of z gives

$$Q(z, \lambda) = \sum_{j=0}^r p_j(z)(\lambda)_j z^{r-j} = 0. \quad (477)$$

The polynomial $Q(z, X)$ is nonzero because $[X^r]Q = p_r(z)$, so λ is algebraic. The coefficient identity and noncancellation are [\[kernel-verified Lean\]](#) in `DfiniteExponentCore`; the differential-equation specialization is the [\[paper proof in this report\]](#) paper layer. This would force integrality for the conditional power branch, whose nonintegral exponent is already known transcendental. Allowing coefficients that depend on x makes differential algebraicity vacuous: $zy' - xy = 0$. Even without admitting x as a coefficient, every power branch satisfies the coefficient-free nonlinear identity

$$zyy'' - z(y')^2 + yy' = 0,$$

so unrestricted nonlinear differential algebraicity still does not isolate the arithmetic exponents.

Integer samples accumulating at a puncture are also insufficient. The entire product below makes $1 + E(1/z)$ holomorphic on $0 < |z| < \varepsilon$, essential at zero, yet equal to 1 at every $z = 2^{-a}3^{-b}$. Exact dilation information, not sampling alone, is indispensable.

88.2 The sharp cubic logarithmic type

Define the locally uniformly convergent entire product

$$E(z) = \prod_{a,b \geq 0} \left(1 - \frac{z}{2^a 3^b}\right). \quad (478)$$

Its zeros are simple and exactly the positive 3-smooth integers. Factorwise equality at $z = -R$ gives $M_E(R) = E(-R)$, and lattice integration yields the exact leading growth

$$\log M_E(R) = \frac{(\log R)^3}{6 \log 2 \log 3} + O((\log R)^2). \quad (479)$$

Indeed, for $T = \log R$,

$$\sum_{a,b \geq 0} (T - a \log 2 - b \log 3)_+ = \frac{T^3}{6 \log 2 \log 3} + O(T^2).$$

For an entire F , set

$$\tau_3(F) = \limsup_{R \rightarrow \infty} \frac{\log^+ M_F(R)}{(\log R)^3}, \quad \tau_* := \frac{1}{6 \log 2 \log 3}.$$

Jensen's formula gives the matching lower bound

$$F \not\equiv 0, \quad F(2^a 3^b) = 0 \ (a, b \geq 0) \implies \tau_3(F) \geq \tau_*. \quad (480)$$

The constant is sharp by (479). These self-contained statements are [\[paper proof in this report\]](#) to this corpus; no claim of priority in the literature is made.

The threshold gives a sharp sufficient analytic interface. More generally, let $p, q > 1$ be multiplicatively independent, $P, Q \in \mathbb{C}^\times$, and suppose an entire H satisfies

$$H(p^a q^b) = P^a Q^b \quad (a, b \geq 0).$$

If

$$\limsup_{R \rightarrow \infty} \frac{\log^+ M_H(R)}{(\log R)^3} < \frac{1}{6 \log p \log q}, \quad (481)$$

then $H(z) = z^n$, $P = p^n$, and $Q = q^n$ for some $n \in \mathbb{N}_0$. To prove this, the defect $H(pz) - PH(z)$ vanishes on the whole lattice and remains subcritical, so (480) annihilates it; coefficient comparison then leaves a single monomial. Thus a subcritical interpolant of the hypothetical data would prove the conjecture. At equality, $1 + E$ is transcendental entire and integer-valued on every node, while $z^n + cE(z)$ interpolates the same integral-exponent data as z^n . These examples show failure of uniqueness at critical type; they do not prove that every more structured critical-type criterion is incapable of forcing exponent integrality.

88.3 The critical boundary cocycle

The product has exact edge factorizations

$$E(2z) = E(z)B_2(z), \quad E(3z) = E(z)B_3(z), \quad (482)$$

where

$$B_2(z) = \prod_{b \geq 0} \left(1 - \frac{2z}{3^b}\right), \quad B_3(z) = \prod_{a \geq 0} \left(1 - \frac{3z}{2^a}\right).$$

Their second-order growth is

$$\log M_{B_2}(R) = \frac{(\log R)^2}{2 \log 3} + O(\log R), \quad \log M_{B_3}(R) = \frac{(\log R)^2}{2 \log 2} + O(\log R). \quad (483)$$

For an entire interpolant $H(2^a 3^b) = M^a A^b$, set $D_2(z) = H(2z) - MH(z)$ and $D_3(z) = H(3z) - AH(z)$. Their lattice zeros and the simplicity of E give unique entire quotients $D_2 = EU$ and $D_3 = EV$. Commutativity of the two dilations is exactly the cocycle

$$B_3(z)U(3z) - AU(z) = B_2(z)V(2z) - MV(z). \quad (484)$$

This is an infinite analytic analogue of `BoundaryDilationDefect`, not a replacement for its finite algebra.

If U and V are polynomials, the strict quadratic gap in (483), evaluated on the negative real axis, first forces $U = 0$ and then $V = 0$. Both dilation equations are exact, so coefficient comparison gives $H(z) = z^n$, $M = 2^n$, and $A = 3^n$. This polynomial-boundary-quotient closure and the cocycle are [\[paper proof in this report\]](#) paper theorems. Critical type alone does not establish the polynomial quotient hypothesis; manufacturing that finite-dimensional boundary regularity is the new open target.

Status: Report 27 was audited as an independent sibling. Its D-finite noncancellation core is [\[kernel-verified Lean\]](#) in `DfiniteExponentCore`. The universal-cover classification, sharp canonical-product/Jensen threshold, subcritical interpolation rigidity, edge growth, critical cocycle, and polynomial-quotient closure are [\[paper proof in this report\]](#) paper results. Overlapping monodromy material is omitted, and neither critical-type sampling nor the cocycle currently supplies the missing interpolant regularity.

A Status matrix for major statements

The status tags have the following meanings. [\[kernel-verified Lean\]](#) marks a statement that the Lean kernel accepts at the current repository state, with no `sorry`, no extra axioms, and no `native_decide`. [\[paper proof in this report\]](#) marks a paper proof given in this report that has not yet been formalized. [\[Lean draft, not yet accepted\]](#) marks a statement for which Lean source exists in the repository but has not yet been accepted by the kernel. [\[software computation\]](#) marks a finite software computation with its local trust boundary—for example CPython/`mpmath` interval arithmetic or C/MPFR/GMP directed rounding—which is not a kernel certificate. [\[classical/open background\]](#) marks classical or open background. No entry below asserts the full Alaoglu–Erdős conjecture; it remains open.

No row of the table carries [\[Lean draft, not yet accepted\]](#) any longer. Every module that earlier revisions of this report listed as a pending draft has since been accepted by the kernel and placed on the aggregate import surface. The tag still applies to the separate, unassembled 16384 finite-certificate attempt described in Section 25, which is not a row of this major-statement table and supplies no theorem.

Statement	Status	Location
<i>Classical frame</i>		
$2^x, 3^x, 5^x \in \mathbb{Z} \Rightarrow x \in \mathbb{Z}$	[kernel-verified Lean]	IntegerExponent
Gelfond–Schneider theorem	[kernel-verified Lean]	GelfondSchneider
Four exponentials conjecture (hypothesis only)	[classical/open background]	[13, 14]
Full Alaoglu–Erdős conjecture	Open	Section 6, Section 27
<i>Transcendence and reduction</i>		
Nonintegral solution irrational and transcendental	[kernel-verified Lean]	TwoBaseIntegerExponent, Transcendence
$\vartheta = \log_2 3$ transcendental	[kernel-verified Lean]	Transcendence
Four-exponentials conjecture $\Rightarrow$ Conjecture 2.1	[kernel-verified Lean]	FourExponentialsReduction
Prime-log-product independence $\Rightarrow$ Conjecture 2.1	[kernel-verified Lean]	Section 12; LogProductIndependence
Prime-log-product independence $\Rightarrow$ the rational-output form (22)	[kernel-verified Lean]	Section 12; LogProductIndependence
Two-prime case of the independence criterion, and its sharpness	[kernel-verified Lean]	Section 12.5
<i>Zero-free regions and finite searches</i>		
$2^x < 100$ and $2^x < 256 \Rightarrow x \in \mathbb{Z}$	[kernel-verified Lean]	FiniteCheck, FiniteCheck256
$2^x < 4096 \Rightarrow x \in \mathbb{Z}$ (so $x \geq 12$)	[kernel-verified Lean]	FiniteCheck4096
No nonintegral candidate with $2^x < 2^{20}$	[software computation]	Section 25, Appendix F
No nonintegral candidate with $2^x < 2^{27}$	[software computation]	Section 40
Reported scan through 2^{29} ; canonical logs and replay artifacts absent from the merged report directory	[software computation]	Section 65.7
<i>Exact conditional structure</i>		
Rank of candidate exponent vectors ≤ 2	[kernel-verified Lean]	MultiplicativeRank
Cross-valuation identity and conditional box growth	[kernel-verified Lean]	CandidateStructure
Common odd core $2^i w^j$ (existence)	[kernel-verified Lean]	OddCore
Counting $\log_2 N \cdot \log_3 N$; linear dyadic blocks	[kernel-verified Lean]	OddCore
Primitive (canonical) odd core; unique coordinates	[kernel-verified Lean]	PrimitiveOddCore
Unique least solution; affine primitive output pair	[kernel-verified Lean]	LeastSolution
Simultaneous primitive output normalization; gcd one; size dichotomy	[kernel-verified Lean]	OutputNormalization
Power-of-three denominator normalization	[kernel-verified Lean]	ThreeDenominatorNormalization
Rational-power index and exact solution monoid $\mathbb{N}_0 + \mathbb{N}_0 \beta$	[kernel-verified Lean]	RationalPowerIndex, PrimitiveGenerator
Canonical primitive generator; exact candidate monoid $\{2^n M^k\}$	[kernel-verified Lean]	PrimitiveGenerator
Denominator-controlled rational third output and determinant margin	[kernel-verified Lean]	RationalThirdBase
Unrestricted rational-third-output determinant grid	[kernel-verified Lean]	RationalSixExponentials, RationalBaseThird

Statement	Status	Location
3-smooth classification of further bases	[kernel-verified Lean]	ThreeSmooth
Both outputs 3-smooth $\Rightarrow x \in \mathbb{Z}$ (unconditional exclusion)	[kernel-verified Lean]	SmoothOutputs
Sharp form: a 3-smooth output forces a partner with a prime factor ≥ 5	[kernel-verified Lean]	SmoothOutputs
Primitive output cores are never exactly 3 and 2; exact core rationality classification	[kernel-verified Lean]	CoreIndependence
<i>Algebraic auxiliary bases and output loci</i>		
Six exponentials for three positive rational bases and rational outputs	[kernel-verified Lean]	RationalSixExponentials
Field-norm determinant bound over rings of integers	[kernel-verified Lean]	AlgebraicIntegerDeterminant
One number field and one natural denominator for a finite algebraic family	[kernel-verified Lean]	AlgebraicOutputBridge, FiniteAlgebraicOutputBridge
Six exponentials for three positive algebraic real bases and algebraic outputs	[kernel-verified Lean]	AlgebraicSixExponentials
Algebraic output at a natural third base	[kernel-verified Lean]	AlgebraicThirdBase
Exact rational-base spectrum: $q^x \in \mathbb{Q}$ iff $x \in \mathbb{Z}$ or q is a 2, 3-unit	[kernel-verified Lean]	RationalBaseThird
Exact algebraic-output form of the same classification	[kernel-verified Lean]	AlgebraicRationalBase
Divisible-hull criterion: an algebraic base with algebraic power has a positive 2, 3-unit power	[kernel-verified Lean]	AlgebraicBaseUnitPower
Pointwise algebraic natural-base classification and iterated-output transcendence	[kernel-verified Lean]	AlgebraicConsequences
Exact algebraic output monomial locus is a coordinate face or the origin	[kernel-verified Lean]	AlgebraicOutputLocus
Algebraic output ratios lie on one external p -adic line (lattice rank ≤ 1)	[kernel-verified Lean]	AlgebraicOutputLocus
Common algebraic-output exponent plane: $y \in \mathbb{Q} + \mathbb{Q}x$	[kernel-verified Lean]	AlgebraicExponentLocus
Exact rational-function locus	[kernel-verified Lean]	RationalFunctionAlgebraicLocus
$P/Q \in \mathbb{Q} + \mathbb{Q}x$ iff $P = (q + rX)Q$; canonical $\mathbb{N}_0$ -cone comparison		
General two-base saturated spectrum	[paper proof in this report]	algebraic spectrum (this report)
$\Gamma(a, b) = \{a^r b^s : r, s \in \mathbb{Q}\}$		
Strong spectrum: $x \log \alpha \notin \mathcal{L}$ for $\alpha \in \overline{\mathbb{Q}}_{>0} \setminus \Gamma$	[paper proof in this report]	algebraic spectrum (this report)
Every $\alpha \in \overline{\mathbb{Q}}_{>0} \setminus \Gamma$ is a one-output detector for Conjecture 2.1	[paper proof in this report]	algebraic spectrum (this report)
<i>One-base equivalent formulations</i>		
Conjecture 2.1 $\iff$ every solution has $5^x \in \overline{\mathbb{Q}}$	[kernel-verified Lean]	AlgebraicConsequences
Conjecture 2.1 $\iff$ every solution has $2^{x^2}, 3^{x^2} \in \overline{\mathbb{Q}}$	[kernel-verified Lean]	AlgebraicConsequences
Conjecture 2.1 $\iff$ every solution has $6^{x^2} \in \overline{\mathbb{Q}}$	[kernel-verified Lean]	AlgebraicConsequences

Statement	Status	Location
Conjecture 2.1 $\iff$ the algebraic locus among $(2^i/3^j)^{x^2}$ has rank two	[kernel-verified Lean]	AlgebraicOutputLocus
Conjecture 2.1 $\iff$ the algebraic real-power output locus has rank three	[kernel-verified Lean]	AlgebraicRankThreeEquivalence
Conjecture 2.1 $\iff$ no squared-monomial algebraic coordinate face	[kernel-verified Lean]	AlgebraicOutputLocus
Conjecture 2.1 $\iff 2^y, 3^y \in \overline{\mathbb{Q}}$ at every natural power $y = x^n$	[kernel-verified Lean]	AlgebraicExponentLocus
Conjecture 2.1 $\iff$ both outputs are algebraic after evaluation of any fixed $P \in \mathbb{Q}[X]$ of degree at least two	[kernel-verified Lean]	RationalFunctionAlgebraicLocus
<i>Canonical radical: valuation, degree, and hull</i>		
Odd-prime valuation gcd; outer-power exponents of the canonical radical	[kernel-verified Lean]	CanonicalRadicalValuation
Exact rational outer-power criterion; the triple gcd is not replaceable by either pairwise gcd	[kernel-verified Lean]	CanonicalRadicalSharpness
Norm and trace of the pure-radical generator; the base-swapped relation blocks monomial injectivity	[kernel-verified Lean]	CanonicalRadicalFieldNorm
Algebraic integrality of the radical iff trivial reduced denominator	[kernel-verified Lean]	CanonicalRadicalThirdOutput
Algebraicity of the radical's β -th power iff 2, 3-smooth numerator	[kernel-verified Lean]	CanonicalRadicalAlgebraicOutput
External numerator prime excludes the positive 2, 3-unit divisible hull	[kernel-verified Lean]	CanonicalRadicalDivisibleHull
Mixed radical monomials retain an external prime; transcendental outputs	[kernel-verified Lean]	CanonicalRadicalMixedOutput
Exact rational-power indices and radical degrees of the simultaneous radicals; product degree de	[kernel-verified Lean]	SimultaneousRadicalDegrees
Localized-radical conjecture target itself	Open	Section 6
<i>Valuation lattice, kernel, and towers</i>		
Rational-output lattice Λ_x of the prime-valuation matrix V_x	[paper proof in this report]	valuation lattice (this report)
Smith classification: $\Lambda_x/\mathbb{Z}^2 \cong \mathbb{Z}/d_1\mathbb{Z} \oplus \mathbb{Z}/d_2\mathbb{Z}$ and the finite rational-output defect	[paper proof in this report]	valuation lattice (this report)
Second-iterate kernel $K_x = \{(r, s) \in \mathbb{Q}^2 : M^r A^s \in \Gamma\}$; $\dim_{\mathbb{Q}} K_x \leq 1$	[kernel-verified Lean]	SecondIterateKernel
Kernel invariance $K_x = K_y$ across nonintegral solutions	[kernel-verified Lean]	SecondIterateKernel
Möbius criterion: $K \neq \{0\}$ iff $x \in \mathrm{PGL}_2(\mathbb{Q}) \cdot \vartheta$, iff $1, \vartheta, x, x\vartheta$ are $\mathbb{Q}$ -dependent	[kernel-verified Lean]	KernelDichotomy
Multiplicative-dependence form: $K \neq \{0\}$ iff $M, A, 2, 3$ are multiplicatively dependent	[kernel-verified Lean]	KernelDichotomy

Statement	Status	Location
Four-type classification of counterexamples by external valuation rank	[paper proof in this report]	second-iterate kernel (this report)
No nontrivial two-step return inside Γ	[paper proof in this report]	depth-three tower (this report)
Depth-three tower rigidity: some α^{x^j} , $j \leq 3$, is transcendental, with exact first depth	[paper proof in this report]	depth-three tower (this report)
Universal three-level detector: Conjecture 2.1 $\iff \alpha^x, \alpha^{x^2}, \alpha^{x^3} \in \overline{\mathbb{Q}}$ for one fixed $\alpha \neq 1$	[paper proof in this report]	depth-three tower (this report)
Base-2 cubic-tower equivalence: Conjecture 2.1 $\iff 2^{x^2}, 2^{x^3} \in \overline{\mathbb{Q}}$	[kernel-verified Lean]	TowerRigidity
Base-2 depth-two dichotomy: $2^{x^2} \in \overline{\mathbb{Q}}$ iff 2^x is 3-smooth, and then $2^{x^2} \in \mathbb{N}$ while 2^{x^3} is transcendental	[kernel-verified Lean]	TowerRigidity
Exact transition $2^y, 2^{xy} \in \overline{\mathbb{Q}}$ iff $y \in \mathbb{Q} + \mathbb{Q}\vartheta$; at most one positive depth starts an adjacent algebraic pair; two fixed distinct depths give an AE equivalence	[kernel-verified Lean]	AlgebraicTowerTransition
Exact base-3 transition locus; no positive base-2/base-3 pair can coexist; across both principal towers at most one positive transition depth, with higher rational-unit directions excluded in the dependent-output branch	[kernel-verified Lean]	CrossBaseTowerTransition
Symmetric tower constants: at least one of $\log E_3, \log E_2$ lies outside $\mathcal{L}$	[paper proof in this report]	symmetric tower constants (this report)
Exact conditional classification of $E_3 = 3^\vartheta$ and $E_2 = 2^{1/\vartheta}$	[kernel-verified Lean]	SymmetricTowerConstants
E_3 and E_2 are never both algebraic under a nonintegral solution	[kernel-verified Lean]	SymmetricTowerConstants
<i>Counting, growth, and equidistribution</i>		
Exact block counts and cumulative finite sum	[kernel-verified Lean]	ExactCounting
Cumulative $T^2/(2\beta) + O(T)$ asymptotic; subquadratic-growth equivalence	[kernel-verified Lean]	QuadraticAsymptotic
Vanishing nonzero block Fourier modes	[kernel-verified Lean]	BlockWeyl
Exact block averages for finite trigonometric polynomials	[kernel-verified Lean]	BlockEquidistribution
Full continuous-test/interval blockwise equidistribution	[paper proof in this report]	Section 5
Density dichotomy, gap equivalences, infinitude	[kernel-verified Lean]	FractionalPartDensity
Nonintegral candidate refinement; density zero	[kernel-verified Lean]	NonintegerDensity
Shrinking-target Diophantine bounds	[kernel-verified Lean]	ShrinkingTarget
Convergent-denominator growth	[kernel-verified Lean]	DenominatorGrowth
$q_{n+1} < 2m^{q_n}$		
<i>Difference and progression obstructions</i>		

Statement	Status	Location
Affine descent, $r + 8k \leq x$	[kernel-verified Lean]	ArithmeticRigidity
Three-term AP obstruction under $h^5 \leq n$	[kernel-verified Lean]	ZeroDensity
Sharp three-term and four-term AP obstructions ($r = 2, 3$)	[kernel-verified Lean]	SharperProgressions, ThirdDifference
Uniform arbitrary-order AP obstruction ($r \geq 2$)	[kernel-verified Lean]	HigherDifference
Effective rational three-term criterion $h^{400} \leq n^{83}$	[kernel-verified Lean]	SharpProgression
<i>Radical and localized reformulations</i>		
Localized-radical and symmetric reformulations	[kernel-verified Lean]	RadicalReformulation, Appendix C
Primitive localized-radical reduction	[kernel-verified Lean]	PrimitiveRadical
Exact radical degree and binomial irreducibility	[kernel-verified Lean]	RadicalDegree
Candidate-form localization equivalence	[kernel-verified Lean]	Localization
<i>New results of this report</i>		
Restricted logarithm-matrix coefficient principle $\iff$ Conjecture 2.1; rational anisotropy of a counterexample	[kernel-verified Lean]	RestrictedMatrixCoefficient; Section 65.1
Sparse irrational-power-curve zero bound, nonzero integral evaluation determinant, and the exact $R + 1$ -term positive-node interpolant	[kernel-verified Lean]	SparsePowerCurve; Section 65.3
Exact rational rectangular orbit-collision count, irrational grid cardinality, and the one-monomial-saving rational-orbit square-grid certificate	[kernel-verified Lean]	RationalOrbitGrid; Section 66.1
Formal prime-symbol determinant and its specialization; formal vanishing exactly in the integral branch; sparse Structural Rank equivalence	[kernel-verified Lean]	FormalPrimeDeterminant; Section 65.2
BK degree threshold $R - 1$ and the factor-four comparison on the $2N$ grid	[paper proof in this report]	Section 65.4
Boundary-dilation defect factorization, unique degree- $K - 1$ residuals, and the degree- $2K - 2$ commutator	[kernel-verified Lean]	BoundaryDilationDefect; Section 65.5
Unbounded $o((\log H)^2)$ and synchronized compact $o(\log H)$ point-counting targets	[paper proof in this report]	Section 65.6
Transcendence of $e^{2\pi i q \beta}$ for $q \in \mathbb{Q}^\times$, deduced from classical six exponentials	[paper proof in this report]	Section 65.6; classical input [classical/open background]
Exponential-polynomial zero count; generalized Vandermonde nonvanishing at ordered nodes	[kernel-verified Lean]	ExponentialPolynomialZeros
Generalized Vandermonde positivity (ordered chamber)	[kernel-verified Lean]	OrderedChamberPositivity
Arbitrary-node determinant integrality and divided-difference repulsion bound	[kernel-verified Lean]	ArbitraryNodeDeterminant

Statement	Status	Location
Universal nonnegative-node superfactorial interpolation divisor, exact consecutive-node witness, and strengthened candidate repulsion/packing	[kernel-verified Lean]	SuperfactorialDeterminant; final repulsion endpoints retain the named mean-value hypothesis
Normalized geometric signed-progression determinant, primitive quotient coprimality, and exact denominator-tax decomposition with bounded remainder	[kernel-verified Lean]	SignedProgressionDeterminant; Section 66.5
Normalized Fejér kernel, exact finite weighted Fourier score, and exact integer-phase hit bound	[kernel-verified Lean]	FejerKernelCertificate; Section 67.2
Uniform metric near-integrality and Haar random-phase extinction Exact denominator-exponent drift, common centered defect, and structural-prime p -adic escape	[paper proof in this report] [kernel-verified Lean]	Section 67.1; finite Fejér bridge separately kernel-checked AdelicDrift; Section 68.2
Sturmian carry telescope, exact prefix mean/variance, joint-carry probability, and covariance kernel	[kernel-verified Lean]	SturmianProbability; Section 68.3
Dense denominator clearing, finite-difference approximation floor, and all-degree normalized tax obstruction	[kernel-verified Lean]	PolynomialApproximationTax; Section 68.4
Rational homogeneous contact nonvanishing and first-order Padé transversality	[kernel-verified Lean]	RationalContactRigidity; Section 68.5
Carry/floor and denominator formal-series identities, including the exact common defect	[kernel-verified Lean]	SturmianGeneratingSeries; Section 69.2
De-integralized common-zero lattice and canonical divisor; compact $\Theta(\log H)$ height dichotomy	[paper proof in this report]	Section 69.1, Section 69.3
Pólya–Carlson natural boundary for the carry, floor, and reduced-denominator series	[classical/open background]	Section 69.2; exact series algebra separately kernel-checked
Fixed 2×2 integral matrix: integral real power iff both original outputs are integral	[kernel-verified Lean]	IntegralPowerMatrix; Section 70.1
Affine-exponent and rational-approximation alternant factorizations	[kernel-verified Lean]	CandidateAlternant; Section 70.3
Cross-Loewner positive integer certificate and optimized two-coset one-half barrier	[paper proof in this report]	Section 70.2, Section 70.3
Fermat-quotient exact identity, power law, first-log determinant interface, and vanishing on $(2^n, 3^n)$	[kernel-verified Lean]	FermatQuotientDeterminant; Section 71.2

Statement	Status	Location
Structural-prime Iwasawa transversality, exact adelic wedge, conditional cascade depth, and Kummer third-order diagnostic	[paper proof in this report]	Section 71; classical Iwasawa kernel where stated
Exact consecutive-logarithmic-mean enclosure	[kernel-verified Lean]	ConsecutiveLogMeanBounds; Section 72.1
Prime-switch antichain, microscopic resonance window, and 2–3 convergent sieve	[paper proof in this report]	Section 72; rational-output switch family only
Exact exclusion of simultaneous 2–3 switches through $b = 10000$	[software computation]	Section 72.3; Python big-integer/rational verifier
Exact two-coordinate grid displacement identities and rank- K boundary defects	[kernel-verified Lean]	BoundaryDisplacement; rank under $K > 0$, nonsingularity, and nonzero multipliers; Section 73.4
Exact rational bilinear minimum, projective arithmetic density, and integral Hadamard-power no-go	[paper proof in this report]	Section 73.1, Section 73.2
Determinantal-divisor compound sandwich and triangular method ceiling	[paper proof in this report]	Section 73.3; ceiling applies to the stated termwise/Frobenius bounds
Two-variable finite-difference divisor for every compound of the square-grid matrix	[paper proof in this report]	Section 73.3; top gain remains lower order
Boundary transfer determinant and Kraus–Loewner positive integral/denominator tax	[paper proof in this report]	Section 73.4, Section 73.5
Integral dilation-difference quotient and its mod-2 derivative reduction	[kernel-verified Lean]	DilationDifferenceQuotient; Section 74.3
Structural-prime denominator lattice, favorable-branch residual scaling, root depths, and odd-degree	[paper proof in this report]	Section 74; only the two displayed residue branches
dilation-resultant valuations		
Prime-power cross-residual cancellation, exact node-divisor scaling, and conditional exact nonzero cross-resultant orders at 2 and 3	[kernel-verified Lean]	CrossResidualScaling, ResultantScaling, Report14NodeDivisorScaling, Report14CrossResultantBridge; primary interpolant scalings and nonzero residue reductions remain hypotheses
Opposite-residual coefficient laws and the full local cross-Sylvester Smith profile in the simultaneous favorable branch	[paper proof in this report]	Section 74.4; branch is nonexhaustive and the residual-specific Smith reduction remains paper-level
Eventual full-content residual profile for $v_2(A), v_3(M) > 0$ and $n > \max(v_2(M), v_3(A))$	[paper proof in this report]	Section 74.4; finite cost and assignment core [kernel-verified Lean] in ContentNormalizedProfileFinite
Convex matching lower bound and $p^{m \binom{m}{2}}$ divisibility for every canonical weighted lower-Toeplitz m -minor	[kernel-verified Lean]	SmithProfileCore
Uniformly negative compound endpoint-rate functional	[kernel-verified Lean]	CompoundEndpointRate; this certifies a no-closure bound, not the conjecture

Statement	Status	Location
Exact structural-prime order of every leading principal subresultant	[paper proof in this report]	Section 74.4; simultaneous favorable coefficient profiles
Sharp direct/shallow displacement-cost inequality and equality classification	[kernel-verified Lean]	DeepSubresultantCost; combinatorial core of the principal-subresultant theorem
Exact $n = 1$ cancellation of the full interpolation denominator from the paired residual resultant	[kernel-verified Lean]	DenominatorCancellationCounterexample; synthetic $(M, A) = (9, 2)$ test data, not a common-exponent pair
Common-grid transformation no-go for simultaneous structural-prime axis support	[kernel-verified Lean]	GridTransformNoGo; nonsingular direction changes cannot enlarge the favorable branch
All- n $\Omega(n^4)$ denominator-cancellation tax for the synthetic pair $(9, 2)$	[paper proof in this report]	Section 74.4; exact paper consequence of the denominator and cross-resultant formulas, not common-exponent data
Universal leading/constant Cramer factors of the paired residual resultant	[paper proof in this report]	Section 74.4; finite $K = 2, 3, 4$ core factorizations are unbundled exact-CAS diagnostics only
Exact unordered-pair decoding from the ordinary and synchronized p -power sums	[kernel-verified Lean]	SynchronizedSumDecoder; Section 75.3
Rational synchronized-sum decoder degree bound $L \leq 2d + 1$	[kernel-verified Lean]	RationalSumDecoder; Section 75.3
Fixed-rank sum entropy, algebraic decoder degree growth, and finite Fourier parity–residue separation	[paper proof in this report]	Section 75; unit-equation input classical where stated
Generic algebraic-base spectrum for an arbitrary independent positive algebraic pair	[kernel-verified Lean]	AlgebraicBaseSpectrum; Section 76.1
Finite rectangular mixed-Mahler inclusion–exclusion identity	[kernel-verified Lean]	MixedMahlerFinite; Section 76.3; no infinite convergence claim
Entire gamma quotient, corrected mixed-Mahler double-pole resonance, and transfer-spectrum diagnosis	[paper proof in this report]	Section 76; holonomic, Pólya–Carlson, Barnes, and q -digamma inputs classical where stated
Exact determinant-one Smith reduction of the geometric Vandermonde, q -Pascal identity, factor formula, divisibility order, and positivity	[kernel-verified Lean]	GeometricVandermondeSmith; Section 77.2
Exterior-minor probability identity, weighted Smith-tail inequalities, and multilevel modular divisor	[paper proof in this report]	Section 77.1, Section 77.2
Rational cross-Loewner Andreief–Selberg determinant, quadratic singular hierarchy, and weighted compound certificate	[paper proof in this report]	Section 77.3; reduced denominator or Smith gain remains open
Infinitely divisible 3-smooth Gram no-go and exact semigroup entropy saturation	[paper proof in this report]	Section 77.4; no closing arithmetic inequality
Finite ordinary two-scale support inclusion–exclusion identity	[kernel-verified Lean]	TwoScaleSupportFinite; Section 78.3; no infinite convergence claim

Statement	Status	Location
Spectral-zeta/logarithmic-Weyl reformulations and complete rank-two heat expansion	[paper proof in this report]	Section 78.1, Section 78.2; exact counting core already kernel-checked, Mellin–Baker expansion paper-level
Sharp fixed-window polynomial node-budget barrier and capacity constant	[paper proof in this report]	Section 78.4; no claim about rational or growing-window approximants
Exact finite geometric Newton expansion, interpolation, omitted tail, and first extrapolation product	[kernel-verified Lean]	GeometricNewtonInterpolation; Section 79.3
Arithmetic unit locking, optimal-block no-go, dyadic certificates, and cofactor minimax unit barrier	[paper proof in this report]	Section 79; no improvement to the accepted finite scan
Rational scalar/matrix-entry first-jet recovery and base-two integral-exponent closure	[kernel-verified Lean]	JordanFirstJet; Section 80.3; matrix real-power formula remains paper-level
Operator local-universality no-go and global/finite/affine/trace/Hankel closing criteria	[paper proof in this report]	Section 80; criteria are not consequences of the two outputs
Exact rational moment-defect denominators and rational two-atom reproduction	[kernel-verified Lean]	Report21MomentArithmetic; Section 82.2
Hausdorff moment representation, finite rational quadrature, and Müntz invisible positive perturbations	[paper proof in this report]	Section 82; full Müntz and Hahn–Banach inputs classical
Hyperbolic rank geometry and rational symmetrization identities	[kernel-verified Lean]	Report22RankGeometry; Section 83
Baker rank-two exclusion, counterexample rank-3/4 classification, and conditional algebraic three-term orbit	[paper proof in this report]	Section 83; Baker input classical
Generic determinant-distance group collision and modular unimodular wrappers	[kernel-verified Lean]	UnimodularGapCollision; Section 84.1
Integer gap gcds, transported rational secants, finite order-signature no-go, and entrywise jet ceiling	[paper proof in this report]	Section 84; jet ceiling is termwise only
Exact rotation-denominator floor law and finite simultaneous worst-case entropy bound	[kernel-verified Lean]	OrbitHeightEntropy; Section 85.1, Section 85.2
Exact compact-orbit heights, Farey limiting distortion, and Baker–contact scale comparison	[paper proof in this report]	Section 85; entropy concerns relabelings, not new graph points
Piecewise projective Padé-defect optimization and efficiency thresholds	[kernel-verified Lean]	Report25ProjectiveDefect; Section 86.1
Number-field projective defect, rational-transfer no-go, and fixed-cardinality group finiteness	[paper proof in this report]	Section 86; external Padé and approximation theorems cited there
Exact finite carry-word telescoping, arity, suffix descent, and ordered reconstruction	[kernel-verified Lean]	Report26CarryPath; Section 87.1

Statement	Status	Location
Multiset carry paths, character balancing, multinomial fibers, and carry-language pressure	[paper proof in this report]	Section 87; exact-rate claims are restricted to the stated carry-language construction
D-finite exponent-relation polynomial and top-coefficient noncancellation	[kernel-verified Lean]	DfiniteExponentCore; Section 88.1
Sharp 3-smooth canonical-product/Jensen threshold, subcritical interpolation rigidity, and critical defect cocycle	[paper proof in this report]	Section 88; polynomial quotient regularity remains an explicit hypothesis
Arbitrary-node repulsion theorem and its packing corollary	[kernel-verified Lean]	ArbitraryNodeRepulsion, conditional on its named divided-difference mean-value hypothesis
Newton divided-difference factorization of the interpolation determinant	[kernel-verified Lean]	ArbitraryNodeRepulsion, unconditional
Semigroup generalized-Vandermonde hierarchy (integrality half)	[kernel-verified Lean]	MixedExponentSemigroup, SemigroupDeterminant
Falling-factorial coefficient bound; dyadic-strip divided-difference bound	[kernel-verified Lean]	FallingFactorialBound
Semigroup gap theorem in quantitative form	[kernel-verified Lean]	SemigroupGapBound
Leibniz determinant bound from separated row/column bounds	[kernel-verified Lean]	SemigroupGapBound, unconditional
Window-span counting bridge for mixed exponents	[kernel-verified Lean]	SemigroupWeightBound
Explicit dyadic $O((\log X)^2)$ bound from that hierarchy	[kernel-verified Lean]	SemigroupGapBound, conditional on the named hypothesis
One-logarithm deficit analysis	[paper proof in this report]	NewtonFactorization
Min-plus ceiling: term-wise p -adic content never exceeds a correct archimedean bound	[kernel-verified Lean]	Section 22
Exact termwise-divisor ceilings (4/5, 7/10, 3/10, $\pi/4$, 0.8710)	[paper proof in this report]	MinPlusCeiling
A unique tropical minimizer forbids cancellation	[kernel-verified Lean]	Section 21
Tropical quotient factorization, all-depth divisibility, and tied-fiber first-layer criterion (no cancellation lower bound)	[kernel-verified Lean]	MinPlusCeiling
Block-preserving tied sums factor into block determinants and consecutive-power Vandermondes; Beatty minimizers, floor gap, quotient divisibility, and exact expansion	[kernel-verified Lean]	TropicalInitialForm
Exact all-layer row-partition expansion and cutoff- K support truncation (no bound on surviving cancellation)	[kernel-verified Lean]	BlockFiberFactorization, TropicalBlockApplication, BeattyTropicalFiber
	[kernel-verified Lean]	RowBlockLaplaceExpansion

Statement	Status	Location
Prefix-crossing lower bound and stars-and-bars count of row partitions below a Beatty excess cutoff (no valuation/cancellation bound)	[kernel-verified Lean]	AssignmentExcessGeometry
Exact Vandermonde collision energy, elementary power-difference valuation, and finite variable-block cube budget	[kernel-verified Lean]	BlockVandermondeValuationBudget
Whole-Vandermonde dyadic and triadic unit fixed divisors under monic column changes (no bound on the residual cascade)	[kernel-verified Lean]	AllLayerUnitBasis, UnitOrderingGain
Uniform scaled-block minor divisors and additive whole-determinant gain $p^{\tau+U} \mid \det$ (no residual-cancellation bound)	[kernel-verified Lean]	RowBlockUnitDivisibility
Column-permuted raw mixed-matrix bridge and additive dyadic/triadic divisors under explicit finite fiber certificates	[kernel-verified Lean]	MixedDeterminantUnitBridge
Genuine global first- N mixed prefix, automatic consecutive fibers, exact dyadic/triadic cutoff expansions, and structural-output divisors (no residual-cancellation upper bound)	[kernel-verified Lean]	GlobalMixedPrefixCascade; strictly decreasing depths, with A odd for the dyadic structural divisor and $3 \nmid M$ for the triadic one
Exact synthetic cascade computations, including counterexamples to first-layer equality and minimum-effective-layer equality	[exact computation]	Section 38.5
Compact synchronized-orbit reformulation	[kernel-verified Lean]	CompactOrbit
Block-growth reformulations (bounded, $o(T)$, $\liminf, \equiv 1$)	[kernel-verified Lean]	BlockGrowthReformulations
Rational-function rigidity, multiplication and quotient criteria, closure reformulations	[paper proof in this report]	Section 8
Perfect-power content, matched affine decompositions, primitive-output statistics	[paper proof in this report]	Section 10
Candidate gcd/lcm lattice; prime-support rigidity of the whole branch	[kernel-verified Lean]	CandidateLattice
Ambient root saturation index d	[paper proof in this report]	Section 11
Operator formulation on the prime-logarithm space	[kernel-verified Lean]	PrimeLogSpaceOperator
Every complex solution is real (principal branch)	[kernel-verified Lean]	ComplexSolutions
Gaussian-integer version is equivalent to the conjecture	[paper proof in this report]	Section 81
Multivalued branches: the log-product wall widened by $\log(-1)$	[paper proof in this report]	Section 81

Statement	Status	Location
Exact cost of the rational antecedent (denominator-support lemma)	[paper proof in this report]	Section 81

B Formalization inventory

All modules live under `Analysis/ExponentialIdentities/Lean/ExponentialIdentities/`, in namespace `LeanProofs`. Every theorem module listed below is reachable from the aggregate import surface; `AuditNewModules` is built separately as an assumption-audit target. The merged tree builds with no errors. None of the theorem modules contains `sorry`, an extra axiom, or `native_decide`; certificate tables are replayed by the kernel.

B.1 Kernel-accepted modules

Module	Contents
<code>IntegerExponent</code> (+3 modules)	three-base theorem, interpolation determinants
<code>TwoBaseIntegerExponent</code>	conjecture statement, candidates, gaps, $2^x < 10$
<code>.../FiniteCheck</code> , <code>FiniteCheck256</code>	zero-free regions $2^x < 100$, $2^x < 256$
<code>.../FiniteCheck4096</code>	zero-free region $2^x < 4096$ (kernel <code>decide</code> table)
<code>.../ZeroDensity</code>	AP obstruction, Roth-number counting, $o(N)$
<code>.../NonintegerDensity</code>	nonintegral candidate refinement, density zero
<code>.../MultiplicativeRank</code>	$\text{rank} \leq 2$, ϑ irrational, $(\log_2 N)^2$
<code>.../CandidateStructure</code>	cross-valuation identity, conditional box growth
<code>.../OddCore</code>	common odd core, $\log_2 N \log_3 N$, dyadic blocks
<code>.../PrimitiveOddCore</code>	gcd-normalized canonical core, unique coordinates
<code>.../LeastSolution</code>	unique positive least solution, exact affine primitivity
<code>.../ThreeDenominatorNormalization</code>	reduced denominator supported at 3
<code>.../RationalPowerIndex</code>	least rational-power index, exact solution monoid
<code>.../PrimitiveGenerator</code>	canonical primitive generator and unique coordinates
<code>.../ExactCounting</code>	exact unit/dyadic block counts and cumulative finite sum
<code>.../QuadraticAsymptotic</code>	sharp quadratic limit; subquadratic-growth equivalence
<code>.../BlockWeyl</code>	exact-block geometric sums and vanishing nonzero Fourier modes
<code>.../BlockEquidistribution</code>	finite trigonometric-polynomial block averages
<code>.../OutputNormalization</code>	simultaneous primitive outputs, gcd one, size dichotomy
<code>.../RationalThirdBase</code>	denominator divisibility and determinant margin
<code>.../RadicalReformulation</code>	localized-radical equivalences, base-swapped form
<code>.../PrimitiveRadical</code>	canonical primitive-power reduction of radical target
<code>.../RadicalDegree</code>	irreducible binomial, exact minimal polynomial and degree
<code>.../SharperProgressions</code>	sharp three-candidate curvature obstruction
<code>.../ThirdDifference</code>	exact third difference, four-candidate obstruction
<code>.../HigherDifference</code>	uniform arbitrary-order forward-difference obstruction
<code>.../PolynomialApproximationTax</code>	exact dense denominator clearing, integer resolution, finite-difference approximation floor, and the all-degree normalized tax obstruction
<code>.../SharpProgression</code>	effective rational three-term criterion $h^{400} \leq n^{83}$
<code>.../ArithmeticRigidity</code>	descent, $r + 8k \leq x$
<code>.../ThreeSmooth</code>	3-smooth classification of further bases
<code>.../ShrinkingTarget</code>	fractional-part shrinking targets, Diophantine lower bounds
<code>.../DenominatorGrowth</code>	convergent-denominator growth cap $q_{n+1} < 2m^{q_n}$
<code>.../Transcendence</code>	Gelfond–Schneider applications: solutions and $\log_2 3$

Module	Contents
.../FractionalPartDensity	density dichotomy, gap equivalences, infinitude
.../Localization	candidate-form localization equivalence
.../FourExponentialsReduction	four exponentials $\Rightarrow$ the conjecture
GelfondSchneider/...	Gelfond–Schneider theorem (port of Karatarakis)
<i>Rational and algebraic auxiliary bases</i>	
.../RationalSixExponentials	six exponentials for three positive rational bases and rational outputs
.../RationalBaseThird	specialization to bases 2, 3; exact 2, 3-unit classification of positive rational third bases
.../AlgebraicOutputBridge	a number field, an embedding, and one natural denominator clearing an algebraic value to an algebraic integer
.../FiniteAlgebraicOutputBridge	one common number field and one denominator for a finite family of algebraic values
.../AlgebraicIntegerDeterminant	field-norm lower bound for a nonzero algebraic-integer determinant; house and factorial bounds on the n^{11} scale
.../AlgebraicSixExponentials	field-norm six-exponentials determinant for three positive algebraic real bases and three algebraic outputs
.../AlgebraicThirdBase	three-base determinant for algebraic output at a natural third base, first over a number field and then by clearing denominators
.../AlgebraicRationalBase	algebraic outputs at positive rational bases; the algebraic locus is exactly the 2, 3-units unless the exponent is integral
.../AlgebraicBaseUnitPower	positive 2, 3-unit powers and the multiplicative rank of an algebraic auxiliary base
.../AlgebraicBaseSpectrum	arbitrary independent positive algebraic pair; exact positive-power/divisible-hull classification of positive algebraic bases with algebraic x th output
.../AlgebraicConsequences	pointwise algebraic natural-base classification; iterated-output transcendence; the 5^x , 6^{x^2} and both-iterate equivalences
.../AlgebraicOutputLocus	exact algebraic output monomial locus (coordinate face or origin); external p -adic line for output ratios; coordinate-face and rank-two equivalences
.../AlgebraicRankThreeEquivalence	rank three in the algebraic real-power output locus forces integrality; the resulting equivalence
.../AlgebraicExponentLocus	the common algebraic-output exponent plane $\mathbb{Q} + \mathbb{Q}x$, and its natural-power equivalence
.../RationalFunctionAlgebraicLocus	exact rational-function affine identity at a counterexample; nonaffine rational-function and degree- ≥ 2 polynomial detectors; canonical rational affine plane versus nonnegative integral cone
<i>Canonical radical arithmetic</i>	
.../CanonicalRadicalValuation	odd-prime valuation gcd; divisors of the gcd are honest outer-power exponents
.../CanonicalRadicalSharpness	exact criterion for a reduced $c/3^a$ to be a rational outer power; the triple gcd resists both pairwise relaxations
.../CanonicalRadicalFieldNorm	norm and trace of the pure-radical generator; the base-swapped relation blocks monomial injectivity
.../CanonicalRadicalThirdOutput	algebraic integrality of the normalized radical is equivalent to a trivial reduced denominator
.../CanonicalRadicalAlgebraicOutput	algebraicity of the radical's β -th power is equivalent to 2, 3-smoothness of the normalized numerator
.../CanonicalRadicalDivisibleHull	an external numerator prime excludes the whole positive 2, 3-unit divisible hull
.../CanonicalRadicalMixedOutput	genuinely mixed radical monomials keep an external prime, so all their outputs are transcendental

Module	Contents
.../SimultaneousRadicalDegrees	the displayed exponents are the least rational-power indices; exact radical degrees and the product degree de
.../CoreIndependence	the two primitive output cores are never exactly 3 and 2; exact rationality classification of the simultaneous cores
<i>Smoothness and log-products</i>	
.../SmoothOutputs	both outputs 3-smooth forces an integral exponent; sharp form supplying a partner with a prime factor ≥ 5
.../LogProductIndependence	prime-log-product independence as a sufficient condition distinct from four exponentials, in both the integer-output and rational-output forms
<i>Determinant hierarchy and reformulations</i>	
.../ArbitraryNodeDeterminant	interpolation determinant at arbitrary ordered nodes; integrality and the divided-difference repulsion bound
.../SuperfactorialDeterminant	universal nonnegative-node superfactorial divisor and exact consecutive-node witness; strengthened candidate determinant, repulsion, and packing bounds
.../SignedProgressionDeterminant	geometric Vandermonde factorization; primitive invariant quotient and coprimality; exact finite logarithmic denominator-tax identity
.../GeometricVandermondeSmith	integral q -Pascal and Newton-basis matrices; determinant-one two-sided reduction to the exact positive Smith diagonal for $q \geq 2$
.../GeometricNewtonInterpolation	exact finite q -Newton power expansion; interpolation through the truncation degree, omitted-tail identity, and first extrapolation product over arbitrary integer parameters
.../FejerKernelCertificate	normalized finite Fejér kernel; exact weighted Fourier score; exact phase and integer-phase hit-count certificates
.../ExponentialPolynomialZeros	zero-count theorem for real exponential polynomials and nonvanishing of the generalized Vandermonde determinant
.../SparsePowerCurve	sparse zero bound on an irrational power curve; nonzero integral evaluation determinant; exact-support positive-node interpolant
.../RationalOrbitGrid	exact rational rectangle collision count; irrational grid cardinality and sparse lower bound; canonical one-monomial-saving rational-orbit certificate
.../RationalContactRigidity	rational polynomial evaluation at a transcendental direction; homogeneous leading-contact nonvanishing and Padé transversality
.../RestrictedMatrixCoefficient	concrete logarithm matrix, rank-one coefficient factorization, rational anisotropy, and exact equivalence with the conjecture
.../IntegralPowerMatrix	explicit real powers of a fixed split integral matrix and their exact entry-integrality/conjecture equivalence
.../JordanFirstJet	rationality recovery from a scalar value and its nonzero first jet; rational matrix-entry extraction; base-two integral-exponent closure conditional on the formal first-jet matrix
.../Report21MomentArithmetic	rational two-atom moment reproduction; four exact common-denominator identities and their structural-prime local-unit conclusions
.../Report22RankGeometry	hyperbolic normal identity, ternary-conic determinant, explicit rational symmetrizer, and symmetric rank-one outer-product normal form
.../UnimodularGapCollision	determinant-distance power collisions in a commutative group and forward/reverse modular and unimodular congruence wrappers

Module	Contents
.../OrbitHeightEntropy	exact integral-slope and logarithmic floor identities; finite injective-index entropy forcing simultaneous two-coordinate denominator bounds
.../Report25ProjectiveDefect	exact piecewise-linear optimization of the abstract weight-two defect and the resulting depth and arithmetic-efficiency thresholds
.../Report26CarryPath	exact distinguished-descendant carry-word sum and length; strict suffix-degree descent and reconstruction from endpoint plus ordered suffix vertices
.../DfiniteExponentCore	bivariate exponent-relation polynomial for a formal linear differential operator; exact top exponent coefficient and nonzero certificate
.../CandidateAlternant	affine real-exponent Vandermonde factorization and the exact rational-approximation candidate alternant
.../SynchronizedSumDecoder	strict monotonicity on fixed-sum pairs and exact unordered-pair recovery from ordinary and synchronized real-power sums
.../RationalSumDecoder	exact $L \leq 2d + 1$ bound for a degree- d rational decoder on the synchronized slice $(2^n + 1, 3^n + 1)$
.../FormalPrimeDeterminant	multivariate prime-symbol determinant, specialization at prime logarithms, exact formal-zero classification, and Structural Rank equivalence
.../FermatQuotientDeterminant	exact Fermat quotient and power law modulo p ; first local-log determinant interface and integral-power determinant vanishing
.../ConsecutiveLogMeanBounds	exact two-sided enclosure for the reciprocal consecutive logarithmic mean, with an explicit $1/(24H^2)$ remainder
.../BoundaryDisplacement	exact two-coordinate evaluation-matrix displacement identities, terminal-boundary projector ranks, and rank- K displacement consequences under explicit $K > 0$, nonsingularity, and nonzero-multiplier hypotheses
.../BoundaryDilationDefect	rectangular node products, dilation-defect divisibility, unique square-grid residuals, and the degree- $2K - 2$ boundary commutator
.../DilationDifferenceQuotient	integral factorization of $P(cX) - P(X)$; dyadic $2X$ quotient and its formal-derivative reduction modulo two
.../CrossResidualScaling	generic prime-power cancellation through a unit polynomial; normalized opposite-residual transfer through a scaled defect identity
.../ResultantScaling	exact resultant change under $X \mapsto sX$; symmetric cross-normalization identities, nonzero residue quotients, and two-prime divisibility
.../Report14NodeDivisorScaling	exact dyadic and triadic normalization of the two report-14 boundary node divisors, with structural-prime-unit constant terms
.../Report14CrossResultantBridge	conditional composition of the report-14 primary and opposite residual scalings into one nonzero cross resultant with exact 2- and 3-orders
.../DenominatorCancellationCounterexample	exact smallest-grid certificate that the full interpolation denominator can disappear from the paired rational residual resultant
.../GridTransformNoGo	exact classification showing that one nonsingular common direction change cannot manufacture simultaneous dyadic and triadic axis separation
.../SmithProfileCore	convex injective-matching cost; strict equality case; universal $p^{m \binom{m}{2}}$ divisibility of weighted lower-Toeplitz minors

Module	Contents
.../ContentNormalizedProfileFinite	endpoint-gap, quotient-cost, direct-floor, wrap-path, and compressed-permutation inequalities for the eventual primitive residual Smith profile
.../DeepSubresultantCost	sharp direct/shallow displacement-cost bound, strict case, and equality classification for the leading principal-subresultant matching
.../CompoundEndpointRate	kernel-checked negativity of the normalized structural compound gain after the mandatory base-2/base-3 endpoint costs
.../OrderedChamberPositivity	positivity of generalized real-power Vandermonde determinants and of the semigroup determinant on ordered chambers
.../MixedExponentSemigroup	pair-indexed exponent $a + b\vartheta$, injectivity, and integrality of $m^{a+b\vartheta}$
.../MixedMahlerFinite	exact finite rectangular inclusion–exclusion identity behind the mixed (2,3)-Mahler equation; no convergence or natural-boundary assertion
.../TwoScaleSupportFinite	exact finite rectangular inclusion–exclusion identity behind the ordinary two-scale support equation; no convergence assertion
.../SemigroupDeterminant	integrality half of the semigroup generalized-Vandermonde hierarchy, with $ \det \geq 1$
.../FallingFactorialBound	falling-factorial coefficient bound and dyadic-strip divided-difference bound
.../SemigroupWeightBound	mixed-exponent family construction and the window-summation counting bridge
.../BlockGrowthReformulations	boundedness, $o(T)$, $\liminf$, and $\equiv 1$ reformulations of the conjecture
.../CompactOrbit	compact synchronized-orbit reformulation with shared floor exponent, in both directions
.../AdelicDrift	exact integer denominator exponents, common centered defect, structural-prime normalized orbit norms, and p -adic escape
.../SturmianProbability	carry/floor telescope, exact prefix law, mean, variance, joint-carry intervals, probability, and covariance kernel
.../SturmianGeneratingSeries	exact carry/floor and denominator formal power-series identities, twice-cleared defect, and cross-base cancellation
.../ArbitraryNodeRepulsion	divided differences, the Newton factorization of the interpolation determinant, and the repulsion and packing bounds conditional on the named mean-value hypothesis
.../SemigroupGapBound	Leibniz determinant bound, semigroup span inequality, and the explicit dyadic $10^4(\log X)^2$ block bound conditional on the named Newton factorization
.../TropicalInitialForm	common-power factorization of a Leibniz determinant; all-depth quotient divisibility and the modulo- p tied-fiber criterion, without a cancellation lower bound
.../BlockFiberFactorization	block-preserving signed Leibniz sums as products of block determinants, including variable-size and consecutive-power Vandermonde forms
.../TropicalBlockApplication	first-extra-prime divisibility iff the product of block determinants, or its consecutive-power Vandermonde specialization, vanishes modulo p ; application hypotheses remain explicit
.../BeattyTropicalFiber	exact repeated-weight minimizers for rank-one assignment costs; Beatty floor gap; normalized quotient divisibility; exact block-determinant expansion
.../RowBlockLaplaceExpansion	exact grouping of all assignment layers by row partitions; normalized and cutoff- K expansions into signed products of square unit minors

Module	Contents
.../AssignmentExcessGeometry	exact prefix-crossing energy; Beatty excess lower bound; stars-and-bars binomial count for row partitions below any fixed cutoff
.../BlockVandermondeValuationBudget	exact collision-depth formula for finite Vandermondes; elementary difference-of-powers valuation and variable-block cubic budgets
.../AllLayerUnitBasis	complete-determinant divisors from monic unitriangular or determinant-one column changes; strengthened positive dyadic-odd and triadic-unit p -ordering weights, without a residual-cascade bound
.../UnitOrderingGain	exact unit block-weight identities and finite square/cube aggregate bounds; no triangular-lattice asymptotic or residual-cascade estimate
.../RowBlockUnitDivisibility	row-scaled consecutive-power matrices; a common p -ordering divisor for every assignment-minor product; determinant-one fiberwise basis change; additive divisor $p^{\tau+U}$ for the full rank-one power determinant
.../MixedDeterminantUnitBridge	finite sorted-box mixed-exponent prefix; exact column-permuted identification with the rank-one scaled model; dyadic and triadic additive divisors under explicit block-fiber certificates
.../GlobalMixedPrefixCascade	global initial-segment certificate, automatic lexicographic/colexicographic fiber equivalences, exact truncated cascade identities, and structural-output determinant divisors under the stated decreasing-depth and residue hypotheses
<i>Second iterates, towers, and the candidate lattice</i>	
.../CandidateLattice	coordinatewise gcd and lcm for $2^s w^t$; unconditional closure of the candidate set under gcd and lcm; prime-support rigidity of the branch
.../KernelDichotomy	multiplicative dependence of $M, A, 2, 3$ is equivalent to a rational linear relation among $1, \vartheta, x, x\vartheta$ and to the Möbius orbit
.../SecondIterateKernel	the second-iterate kernel as a subgroup of $\mathbb{Z}^2$; no same-sign vectors, rank at most one, smooth coordinate cases, and solution-independence
.../TowerRigidity	base-two second iterate as a natural number; the depth-two algebraic dichotomy; the depth-three transcendence and the cubic-tower equivalence
.../AlgebraicTowerTransition	exact transition locus $2^y, 2^{xy} \in \overline{\mathbb{Q}} \iff y \in \mathbb{Q} + \mathbb{Q}\vartheta$ under a counterexample; global uniqueness of a positive transition depth, the all-depth three-level exclusion, and the fixed-distinct-depth AE equivalence
.../CrossBaseTowerTransition	exact scaled base-three transition locus; cross-base mutual exclusion, combined uniqueness across both principal towers, and higher smooth-direction exclusions in the dependent-output branch
.../SymmetricTowerConstants	$E_3 = 3^\vartheta$ and $E_2 = 2^{1/\vartheta}$, their conditional rational powers, and their mutual exclusion
.../PrimeLogSpaceOperator	the ϑ -pullback of the prime-logarithm span is exactly $\mathbb{Q} \log 2$ if and only if the rational-output conjecture holds
.../ComplexSolutions	on the principal branch every complex solution is real, by incommensurability of the two argument gratings
.../MinPlusCeiling	the $\{2, 3\}$ -part of an integer divides it, hence term-wise p -adic content can never exceed a correct archimedean bound

Module	Contents
.../AuditNewModules	<code>#print axioms</code> audit of the headline algebraic-locus, rational-function, tower, radical-degree, mixed-output, determinant, tropical, and block-fiber theorems

B.2 Former theorem drafts and the remaining finite-certificate draft

No theorem module in the nine-item list from earlier revisions remains pending. Those revisions listed the following Lean sources as not yet accepted:

ArbitraryNodeDeterminant, ExponentialPolynomialZeros, MixedExponentSemigroup,
SemigroupDeterminant, FallingFactorialBound, SemigroupWeightBound,
BlockGrowthReformulations, CompactOrbit, LogProductIndependence.

All nine have since been accepted by the kernel, placed on the aggregate import surface, and moved into the table above. Statements that those modules actually contain are therefore [\[kernel-verified Lean\]](#).

Separately, the attempted $2^x < 16384$ finite-certificate assembly described in Section 25 remains a genuine [\[Lean draft, not yet accepted\]](#): its individual arithmetic rows replay, but its covering theorem does not compile end to end. It is not on the aggregate import surface and raises no kernel-verified bound. The accepted finite frontier remains $2^x < 4096$.

A second round has since added eight further modules: `CandidateLattice`, `KernelDichotomy`, `SecondIterateKernel`, `TowerRigidity`, `SymmetricTowerConstants`, `PrimeLogSpaceOperator`, `ArbitraryNodeRepulsion` and `SemigroupGapBound`, all on the aggregate import surface and all audited with `#print axioms`.

The current wave adds the following modules:

AlgebraicExponentLocus, AlgebraicOutputLocus,
RationalFunctionAlgebraicLocus, AlgebraicTowerTransition,
CrossBaseTowerTransition, BlockVandermondeValuationBudget, AllLayerUnitBasis,
UnitOrderingGain,
CanonicalRadicalMixedOutput, SimultaneousRadicalDegrees,
TropicalInitialForm, BlockFiberFactorization, TropicalBlockApplication,
BeattyTropicalFiber,
RowBlockLaplaceExpansion, AssignmentExcessGeometry, RowBlockUnitDivisibility,
MixedDeterminantUnitBridge,
RestrictedMatrixCoefficient, FormalPrimeDeterminant, SparsePowerCurve,
BoundaryDilationDefect, SuperfactorialDeterminant, RationalOrbitGrid,
SignedProgressionDeterminant, FejerKernelCertificate,
AdelicDrift, SturmianProbability, SturmianGeneratingSeries,
PolynomialApproximationTax,
RationalContactRigidity, IntegralPowerMatrix, CandidateAlternant,
FermatQuotientDeterminant, ConsecutiveLogMeanBounds, BoundaryDisplacement,
GlobalMixedPrefixCascade, DilationDifferenceQuotient, SynchronizedSumDecoder,
RationalSumDecoder,

AlgebraicBaseSpectrum, MixedMahlerFinite, GeometricVandermondeSmith,
 TwoScaleSupportFinite, GeometricNewtonInterpolation, JordanFirstJet,
 Report21MomentArithmetic, Report22RankGeometry, UnimodularGapCollision,
 OrbitHeightEntropy, Report25ProjectiveDefect, Report26CarryPath, DfiniteExponentCore,
 CrossResidualScaling, ResultantScaling, Report14NodeDivisorScaling,
 Report14CrossResultantBridge,
 DenominatorCancellationCounterexample, GridTransformNoGo, SmithProfileCore,
 ContentNormalizedProfileFinite, DeepSubresultantCost, CompoundEndpointRate.

Their headline APIs are likewise on the aggregate surface and covered by the axiom-audit module.

Two cautions survive both promotions, and the status matrix reflects them. First, a module being accepted certifies exactly the theorems it states. The analytic halves of the repulsion and gap theorems are now formal, but each carries its missing analytic input as an explicit, named *hypothesis* — `RpowDividedDifferenceMeanValue` and `NewtonFactorization` — rather than as an axiom, so those two conclusions are conditional in a way the reader can see in the theorem statement; everything upstream and downstream of them is unconditional. `ExponentialPolynomialZeros` supplies the nonvanishing engine, while the separate accepted module `OrderedChamberPositivity` supplies the ordered-chamber sign proved on paper in Appendix D. Second, the interval scan through 2^{20} and the combined directed-rounding scan through 2^{27} are [\[software computation\]](#), not kernel certificates; the kernel-verified exponent bound is still $x \geq 12$ from `FiniteCheck4096`.

Each further headline theorem should be added to the aggregate import surface and checked with `#print axioms` before its status is promoted. Numerical constants in the tables of Appendix E are explanatory only; the formal statements should use exact expressions or safe rational constants such as $1/24$ and 10000 .

C A symmetric formulation

Everything above has a base-swapped version. Put $\eta = \log_3 2 = 1/\vartheta$. Failure of the conjecture is also equivalent to the existence of an integer $v > 1$ with $3 \nmid v$ and

$$v^\eta \in \mathbb{Z}[1/2].$$

Status: [\[kernel-verified Lean\]](#) —

`not_alaogluErdosConjecture_iff_exists_threeFree_localized_radical` and
`alaogluErdosConjecture_iff_threeFreeLocalizedRadicalExclusion` in
`RadicalReformulation`.

For the primitive pair (M, A) , factoring A into a power of 3 times a 3-coprime part produces a symmetric primitive core and rational-power index; the two normalizations give the same least exponent β , and their compatibility condition is exactly $\log M \log 3 = \log A \log 2$. The symmetry is useful in formalization — some divisibility arguments are easier on one side — and confirms that the localized-radical reduction of Section 6 is not an artifact of privileging the base 2.

D Detailed proof of the generalized Vandermonde sign

This appendix expands the generalized Vandermonde positivity lemma of Section 19 to make its analytic input fully transparent.

Let $0 \leq \lambda_0 < \dots < \lambda_r$ and define $f_j(x) = x^{\lambda_j}$ on $(0, \infty)$. For $k \leq r$,

$$f_j^{(k)}(x) = (\lambda_j)_{\underline{k}} x^{\lambda_j - k}.$$

The k th initial Wronskian is

$$\begin{aligned} W_k(x) &= \det(f_j^{(i)}(x))_{0 \leq i, j \leq k} \\ &= x^{\sum_{j=0}^k \lambda_j - k(k+1)/2} \det((\lambda_j)_{\underline{i}})_{0 \leq i, j \leq k}. \end{aligned}$$

The polynomial $(z)_{\underline{i}}$ is monic of degree i . Replacing the monomial basis $1, z, \dots, z^k$ by the falling-factorial basis is unit lower-triangular. Therefore

$$\det((\lambda_j)_{\underline{i}})_{i,j=0}^k = \prod_{0 \leq i < j \leq k} (\lambda_j - \lambda_i) > 0.$$

Hence every initial Wronskian is positive.

One standard characterization of extended complete Chebyshev systems [6] now implies that for $x_0 < \dots < x_r$,

$$\det(f_j(x_i))_{i,j=0}^r > 0.$$

To avoid using that characterization as a black box, proceed inductively. Suppose a nontrivial combination

$$F(x) = \sum_{j=0}^r c_j x^{\lambda_j}$$

has $r + 1$ distinct positive zeros. Put $x = e^t$ and factor the lowest exponential:

$$e^{-\lambda_0 t} F(e^t) = c_0 + \sum_{j=1}^r c_j e^{(\lambda_j - \lambda_0)t}.$$

Rolle's theorem gives at least r zeros of its derivative, an exponential polynomial with at most r terms. Induction gives a contradiction unless all coefficients vanish. Thus the evaluation determinant is nonzero. Its sign cannot change on the connected chamber $x_0 < \dots < x_r$. Let the nodes coalesce as $x_i = x + i\varepsilon$ and divide by the positive ordinary Vandermonde. As $\varepsilon \downarrow 0$, the quotient tends to $W_r(x) / \prod_{k=0}^r k! > 0$. Hence the original determinant is positive everywhere in the ordered chamber.

Status: [\[paper proof in this report\]](#) and [\[kernel-verified Lean\]](#) — paper proof in this report; kernel-accepted zero-count/nonvanishing in `ExponentialPolynomialZeros` and the positive ordered-chamber determinant in `OrderedChamberPositivity`.

E Numerical constants

For reference,

$$\vartheta = 1.58496250072115618145\dots, \quad \vartheta(\vartheta - 1) = 0.92714362797110482756\dots$$

The three-point span constant is

$$\left(\frac{8}{\vartheta(\vartheta - 1)}\right)^{1/3} = 2.051072395661874\dots$$

For the ordinary-power order r , the raw span constant and exponent are

$$C_r = \left(\frac{r!}{|(\vartheta)_{\underline{r}}|}\right)^{2/(r(r+1))}, \quad \alpha_r = \frac{2(r - \vartheta)}{r(r + 1)}.$$

The continuous critical point for α_r solves

$$-r^2 + 2\vartheta r + \vartheta = 0,$$

so

$$r_* = \vartheta + \sqrt{\vartheta^2 + \vartheta} = 3.6090841940\dots,$$

confirming that $r = 4$ is the optimal integer order.

For the semigroup hierarchy of Section 20, the first exponents of $\Lambda_\vartheta = \{a + b\vartheta : a, b \in \mathbb{N}_0\}$ and their pair labels are:

Index	Pair (a, b)	Exponent $a + b\vartheta$
0	(0, 0)	0
1	(1, 0)	1
2	(0, 1)	ϑ
3	(2, 0)	2
4	(1, 1)	$1 + \vartheta$
5	(3, 0)	3
6	(0, 2)	2ϑ
7	(2, 1)	$2 + \vartheta$
8	(4, 0)	4
9	(1, 2)	$1 + 2\vartheta$
10	(3, 1)	$3 + \vartheta$
11	(0, 3)	3ϑ
12	(5, 0)	5

These decimal values are explanatory. They are not part of any machine-checked statement, and formal versions of the corresponding theorems should carry exact expressions instead.

F Interval computation: script and outputs

The script below verifies one range. Set **START** to 1 or 2^{19} and **LIMIT** to 2^{19} or 2^{20} to obtain the two halves reported afterwards.

```

1 import math, time, hashlib
2 import mpmath as mp
3
4 START = 1
5 LIMIT = 1 << 19
6 DPS = 40
7
8 mp.iv.dps = DPS
9 ln2 = mp.iv.log(2)
10 ln3 = mp.iv.log(3)
11 theta_float = math.log(3.0) / math.log(2.0)
12
13 checked = 0
14 failures = []
15 min_lower = (float("inf"), None, None)
16 min_upper = (float("inf"), None, None)
17 manifest = hashlib.sha256()
18 start_time = time.time()
19
20 for m in range(START, LIMIT):
21     if (m & (m - 1)) == 0:      # integral exponent: m is a power of two
22         continue
23
24     # Only proposes a neighboring integer. The interval tests below certify it.
25     n = math.floor(math.exp(theta_float * math.log(m)))
26
27     lhs = mp.iv.log(m) * ln3
28     lower_diff = lhs - mp.iv.log(n) * ln2
29     upper_diff = mp.iv.log(n + 1) * ln2 - lhs
30
31     lower_lb = float(lower_diff.a)
32     upper_lb = float(upper_diff.a)
33     if not (lower_lb > 0.0 and upper_lb > 0.0):
34         failures.append((m, n, str(lower_diff), str(upper_diff)))
35         break
36
37     min_lower = min(min_lower, (lower_lb, m, n))
38     min_upper = min(min_upper, (upper_lb, m, n))
39     manifest.update(f"{m}:{n};".encode())
40     checked += 1
41
42 print("checked_nonpowers=", checked)
43 print("failures=", len(failures))
44 print("min_lower_log_margin=", min_lower)
45 print("min_upper_log_margin=", min_upper)
46 print("sha256_m_n_manifest=", manifest.hexdigest())
47 print("elapsed_seconds=", time.time() - start_time)
48 assert not failures

```

Listing 12: Interval verification

The two 40-digit runs produced:

```

LIMIT=524288
DPS=40
checked_nonpowers=524268
failures=0
min_lower_log_margin=(1.4147089634746718e-15, 189427, 231498127)

```



```
min_upper_log_margin=(3.0841010256409534e-15, 306982, 497586055)
sha256_m_n_manifest=02b589461fa8c1721bc251575b98b1e3d58001f54cb7faa2cd6c59edb8c819a5
elapsed_seconds=18.505
```

Listing 13: Lower half output.

```
LIMIT=1048576
DPS=40
checked_nonpowers=524287
failures=0
min_lower_log_margin=(9.490232037370244e-16, 958460, 3023921230)
min_upper_log_margin=(2.3019781003173366e-15, 556061, 1275861767)
sha256_m_n_manifest=fcccf5de8673379d37ff0fc840ab091d9ff7a2352ff64a93091f01c21bf4cb4
elapsed_seconds=18.904
```

Listing 14: Upper half output.

The two SHA-256 manifests are

lower: 02b589461fa8c1721bc251575b98b1e3d58001f54cb7faa2cd6c59edb8c819a5,

upper: fcccf5de8673379d37ff0fc840ab091d9ff7a2352ff64a93091f01c21bf4cb4.

Elapsed times are machine-dependent sanity checks; the counts, zero-failure results, extremal tuples, and hashes are the reproducibility data. The hashes cover the ASCII stream `m:n`; of proposed-and-certified pairs and make accidental changes to the candidate sequence detectable across reruns.

Status: [\[software computation\]](#) — interval arithmetic in `mpmath` [15] at 40 digits. This is a software trust boundary, not a kernel certificate. The kernel-verified zero-free region remains $2^x < 4096$ in `FiniteCheck4096`.

G Suggested experimental determinant families

Before developing a full p -adic theory, several finite experiments can test whether the required amplification exists.

G.1 Rectangular row and column boxes

Choose rows (n, k) with $0 \leq n < N_0$, $0 \leq k < K_0$ and columns (a, b) with $0 \leq a < A_0$, $0 \leq b < B_0$. The entry is

$$(2^a 3^b)^n (M^a A^b)^k.$$

For complete rectangles, the matrix is a rearranged Kronecker product only when the row and column cardinalities match in a compatible way. Its determinant or maximal minors can be factored symbolically for small boxes. Record valuations at 2 and 3 as polynomials in $a_0 = v_2(M)$ and $b_0 = v_3(A)$.

G.2 Triangular boxes adapted to height

Use

$$n + k\beta \leq T, \quad a + b\vartheta \leq L.$$

These are the natural row and column shapes from the exact counting theorem and the mixed-exponent semigroup. Search for square minors with minimal archimedean weight and maximal forced valuation. The linear boundary slopes are irrational, so floor-sum fluctuations may prevent cancellation patterns present in rectangles.

G.3 Finite-difference row operations

Order rows by n for fixed k , or by k for fixed n , and apply repeated difference operators. Differences introduce factors such as

$$2^a 3^b - 1, \quad M^a A^b - 1.$$

These may have systematic 2- or 3-adic valuations for carefully selected column pairs. Lifting-the-exponent lemmas can then provide exact lower bounds. The risk is that the archimedean size of the same factors grows too quickly; experiments should track both sides.

G.4 Confluent columns

Instead of only distinct mixed monomials, consider derivatives with respect to an auxiliary continuous exponent parameter before specializing. Confluent generalized Vandermonde matrices introduce logarithmic factors, which are not integral, but common-denominator or linear-form constructions may restore arithmetic control. This is closer to Laurent’s interpolation determinant method [10] and may reduce total exponent weight.

H Dependency map and Lean-oriented proof sketches

The conditional arithmetic criteria of this report were written with formalization in mind: each one is a short argument over an already kernel-verified base. This appendix records the shortest route from the existing corpus to each of them, and then gives skeleton statements for the four criteria whose formalization is least obvious.

Status: Design note. The snippets below are schematic ASCII transliterations and are not claimed to compile unchanged. The mathematics they target is [\[paper proof in this report\]](#); none of it has been accepted by the kernel, and nothing here may be read as a claim of verification.

H.1 Dependency map

Each row names one new result, the principal existing input it is built on, and the extra mathematics a formalization would have to supply. All of the new results are conditional: they describe the structure forced on $\mathcal{S}$ and $\mathcal{C}$ by a hypothetical counterexample to Conjecture 2.1, and they are vacuous if the conjecture holds.

New result	Principal existing input	Additional mathematics
rational-function rigidity	canonical primitive generator; transcendence of β	injectivity of polynomial evaluation at a transcendental element
product and power criteria	unique (n, k) coordinates	quadratic coefficient comparison
quotient criterion	unique (n, k) coordinates	denominator clearing and coefficient comparison
three-smooth output test	transcendence of ϑ	unique factorization in $\mathbb{N}$
product transcendence	three-smooth-base theorem	classical six exponentials
common power degree	unique (n, k) coordinates	positive root uniqueness; divisibility
primitive statistics	total lattice count	Möbius inversion
gcd and lcm lattice	exact candidate monoid	prime-valuation min and max
ambient root saturation	primitive valuation gcd of w	Bézout and linear cone arithmetic
optimized differences	arbitrary-order difference theorem	gamma identity and Stirling estimate

Every entry in the middle column is [\[kernel-verified Lean\]](#) in the corpus catalogued in Section 4. In the right-hand column, only the six exponentials theorem is external background ([\[classical/open background\]](#)); the remaining items are ordinary mathlib-level algebra and analysis, which is what makes these criteria cheap to formalize once the base is in place.

Transliteration conventions. The listings write Unicode in plain ASCII: $\mathbb{R}$, $\mathbb{Q}$, $\mathbb{Z}$ and $\mathbb{N}$ are the real, rational, integer and natural types; arrows are $\rightarrow$ and $\leftarrow$; membership is `in` and divisibility is `dvd`; anonymous constructors use angle brackets; the two logical connectives use their two-character ASCII spellings; and the focusing dot is written as a period. Coercions are written as explicit casts such as `Int.cast` rather than as the coercion arrow.

H.2 Multiplication skeleton

A direct statement can use the repository predicate `TwoBaseIntegralSolution`, which is the defining conjunction $2^x \in \mathbb{Z}$ and $3^x \in \mathbb{Z}$.

```

theorem product_solution_iff_integer_factor
  {x y : ℝ}
  (hx : TwoBaseIntegralSolution x)
  (hy : TwoBaseIntegralSolution y) :
  TwoBaseIntegralSolution (x * y) <->
    x in Set.range (Int.cast : ℤ → ℝ) \ /
    y in Set.range (Int.cast : ℤ → ℝ) := by
classical
constructor
. intro hxy
by_cases hAE : AlaogluErdosConjecture
. left
exact hAE x hx
. obtain <w, d, a, c, beta, ..., hcoord, ...> :=
exists_canonical_primitiveGenerator_of_not_alaogluErdosConjecture hAE

```

```

obtain <nx, kx, hxcoord> := (hcoord x).mp hx
obtain <ny, ky, hycoord> := (hcoord y).mp hy
obtain <np, kp, hpcoord> := (hcoord (x * y)).mp hxy
-- If kx and ky are both positive, expanding the product gives a
-- nonzero quadratic polynomial over Q vanishing at beta, which
-- contradicts transcendence of beta. Hence kx = 0 or ky = 0.
...
. rintro (hxint | hyint)
. exact solution_nsmul_left hxint hy
. exact solution_nsmul_right hyint hx

```

Listing 15: Product criterion: skeleton.

In the finished theorem the integer range can be narrowed to $\mathbb{N}_0$, because every solution is nonnegative by `nonneg_of_two_rpow_integer`. The algebraic step should be isolated, since three of the four criteria reduce to it. Note that the coefficients must be rational — with arbitrary real a, b, c the statement is false — and that is exactly how the applications supply them, as differences of the integer coordinates:

```

lemma transcendental_quadratic_coeff_eq_zero
  {beta : R} {a b c : Q}
  (hbeta : Transcendental Q beta)
  (h : (a : R) * beta ^ 2 + (b : R) * beta + (c : R) = 0) :
  a = 0 /\ b = 0 /\ c = 0 := ...

```

Listing 16: The shared coefficient lemma.

Mathlib’s polynomial evaluation API is probably the cleanest route: build the polynomial $aX^2 + bX + c$, rewrite `h` as an evaluation, and apply the definition of transcendence to conclude that the polynomial is zero.

H.3 Output smoothness skeleton

Define the predicate

$$\text{ThreeSmooth}(m) :\iff \exists u, v \in \mathbb{N}_0, m = 2^u 3^v.$$

A practical implementation should avoid a partial “natural output” function and instead carry explicit witnesses $M, A \in \mathbb{N}$ for the two integer powers:

```

theorem integer_of_outputs_threeSmooth
  {x : R} {M A : N}
  (hM : (2 : R) ^ x = (M : R))
  (hA : (3 : R) ^ x = (A : R))
  (hMs : ThreeSmooth M) (hAs : ThreeSmooth A) :
  x in Set.range (Int.cast : Z -> R) := by
obtain <a, b, rfl> := hMs
obtain <c, d, rfl> := hAs
-- Taking logarithms base two: x = a + b * theta and x * theta = c + d * theta.
have hpoly :
  (b : R) * logThreeDivLogTwo ^ 2
  + ((a : R) - (d : R)) * logThreeDivLogTwo - (c : R) = 0 := by
...
have htheta := transcendental_logThreeDivLogTwo
-- The coefficients are integers, so all three vanish:
-- b = 0, a = d, c = 0, whence 2 ^ x = 2 ^ a and x = a.
...

```

Listing 17: Output smoothness: skeleton.

The converse direction is immediate: for integral $x = a$ the outputs are 2^a and 3^a . The interesting content is that three-smoothness of *both* outputs is already enough to force integrality, with no appeal to the conditional structure theory.

H.4 Common-power skeleton

The formal theorem is most reusable when it is stated for explicit output witnesses and for a fixed canonical generator, so that the coordinate comparison is available as a hypothesis rather than reconstructed inside the proof.

```
theorem simultaneous_perfectPower_iff_divides_coordinates
  (hfail : not AlaogluErdosConjecture)
  {x beta : R} {n k r : N}
  (hr : 0 < r)
  (hbeta : CanonicalPrimitiveGenerator beta)
  (hxcoord : x = (n : R) + (k : R) * beta) :
  (Exists fun U : N => (U : R) ^ r = (2 : R) ^ x) /\
  (Exists fun V : N => (V : R) ^ r = (3 : R) ^ x) <=>
  r dvd n /\ r dvd k := by
-- Convert the two perfect powers into a solution at x / r,
-- apply uniqueness of coordinates, and compare.
...
```

Listing 18: Common powers: skeleton.

The positive-root step can be discharged either by strict monotonicity of $t \mapsto t^r$ on the nonnegative reals or by injectivity of the logarithm; the first is usually shorter in mathlib.

H.5 Ambient saturation skeleton

The only genuinely arithmetic lemma in the saturation argument is the reconstruction of a root of a power of a primitive number:

```
lemma primitive_core_root
  {w u q r : N}
  (hw : oddPrimeValuationGCD w = 1)
  (hr : 0 < r)
  (hpow : u ^ r = w ^ q) :
  Exists fun t : N => q = r * t /\ u = w ^ t := by
-- Compare every coordinate of the prime factorizations.
-- A Bezout combination of the positive valuations of w gives r | q;
-- factorization extensionality then gives u = w ^ t.
...
```

Listing 19: Primitive-core root lemma.

The proof in this report only needs the case $q = dk$, but the general lemma is worth stating once: the same reconstruction is used implicitly in `PrimitiveOddCore` and in `OutputNormalization`, and a shared version would remove that duplication.

H.6 Suggested theorem order for a pull request

A low-conflict sequence, each step reviewable on its own, is:

1. add a small `ThreeSmooth` API together with the output-smoothness theorem;
2. add coordinate comparison lemmas for affine, quadratic and quotient expressions;
3. prove the product, power and quotient criteria;
4. prove the common-power and affine-descent criteria;
5. add the candidate divisibility and support results;
6. add primitive-core root reconstruction and ambient saturation;
7. only then add the asymptotic and six-exponentials extensions.

This ordering keeps the first several commits elementary, and it defers everything that depends on external transcendence input to the end, where it can be reviewed against the status discipline of Appendix A rather than mixed into routine algebra.

I Directed-rounding verifier

Appendix F gives the interval-arithmetic script behind the 2^{20} verification. This appendix gives the second, independent verifier: a C program that uses MPFR directed rounding rather than an interval library, and that extends the scanned range sixty-fourfold, to 2^{26} .

Status: [software computation] — MPFR 4.2.2 at 192 bits with directed rounding modes. Like the 2^{20} scan this is a software trust boundary, not a kernel certificate. The kernel-verified zero-free region remains $2^x < 4096$ in `FiniteCheck4096`.

I.1 What the verifier certifies

For a non-power of two m , proving $m^\vartheta \notin \mathbb{N}$ amounts to exhibiting $n \in \mathbb{N}$ with $n < m^\vartheta < n + 1$, and, as in Section 25, that pair of strict inequalities is certified in logarithmic form so that no exponentiation is ever performed. At precision $p = 192$ the program computes directed enclosures

$$L_m^- \leq \log m \leq L_m^+, \quad L_2^- \leq \log 2 \leq L_2^+, \quad L_3^- \leq \log 3 \leq L_3^+,$$

together with $L_n^+ \geq \log n$ and $L_{n+1}^- \leq \log(n + 1)$, and then rounds *downward* in

$$\delta_- = L_m^- L_3^- - L_n^+ L_2^+, \quad \delta_+ = L_{n+1}^- L_2^- - L_m^+ L_3^+.$$

Positivity of δ_- proves the lower inequality and positivity of δ_+ proves the upper one. Every input on the n side is below 2^{42} and is therefore exactly representable at 192 bits.

An ordinary `long double` evaluation proposes $n \approx \lfloor m^\vartheta \rfloor$ and the offsets $0, -1, 1, -2, 2, -3, 3$ are tried in that order. The proposal has no logical role: a proposal is accepted only when both directed

lower bounds are strictly positive, and a failure to locate any admissible n is reported as a failure rather than silently reclassified. Locator error can therefore produce a false alarm but not a false proof. Powers of two are skipped rather than tested, since $(2^q)^{\vartheta} = 3^q$ is the known integral candidate; an exact integral value at a non-power of two would satisfy neither strict inequality for any n and would be reported.

I.2 Source identity and environment

The recorded run used MPFR 4.2.2 linked as `libmpfr.so.6`, GCC 14.2.0 with `-O3 -fopenmp`, GNU OpenMP with five worker threads, Linux on x86-64, and 192-bit precision with the directed modes `MPFR_RNDD` and `MPFR_RNDU`. The SHA-256 digest of the exact executed source reproduced below is

f888c0d89be9366f01e2d26f339dcbe446556d39d4c5dbcf8f579ec0d0052e1d

The execution image contained the MPFR shared library but not its development header, so the source carries a guarded minimal declaration of the MPFR 4.x ABI; on a system with `mpfr.h` present the standard header branch is taken instead. That ABI detail is part of the software trust boundary and is disclosed here rather than hidden.

I.3 Source

The listing below is the executed source verbatim. Its language is C99; it is typeset in this report’s generic code style, which has no C-specific highlighting.

```

1 #define _GNU_SOURCE
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <stdint.h>
5 #include <inttypes.h>
6 #include <math.h>
7 #include <omp.h>
8
9 #if defined(__has_include)
10 #   if __has_include(<mpfr.h>)
11 #       include <mpfr.h>
12 #       define HAVE_SYSTEM_MPFR_HEADER 1
13 #   endif
14 #endif
15
16 #ifndef HAVE_SYSTEM_MPFR_HEADER
17 /* Minimal declarations for the MPFR 4.x ABI. The execution environment had
18    libmpfr.so.6 but not the development header package. A normal build with
19    libmpfr-dev installed takes the branch above instead. */
20 typedef unsigned long mp_limb_t;
21 typedef long mpfr_prec_t;
22 typedef long mpfr_exp_t;
23 typedef int mpfr_sign_t;
24 typedef struct {
25     mpfr_prec_t _mpfr_prec;
26     mpfr_sign_t _mpfr_sign;
27     mpfr_exp_t _mpfr_exp;
28     mp_limb_t *_mpfr_d;
29 } __mpfr_struct;

```

```

30 typedef __mpfr_struct mpfr_t[1];
31 typedef __mpfr_struct *mpfr_ptr;
32 typedef const __mpfr_struct *mpfr_srcptr;
33 typedef enum {
34     MPFR_RNDN = 0, MPFR_RNDZ = 1, MPFR_RNDU = 2,
35     MPFR_RNDD = 3, MPFR_RNDA = 4, MPFR_RNDF = 5,
36     MPFR_RNDNA = -1
37 } mpfr_rnd_t;
38 extern void mpfr_init2(mpfr_ptr, mpfr_prec_t);
39 extern void mpfr_clear(mpfr_ptr);
40 extern int mpfr_set_ui(mpfr_ptr, unsigned long, mpfr_rnd_t);
41 extern int mpfr_log(mpfr_ptr, mpfr_srcptr, mpfr_rnd_t);
42 extern int mpfr_mul(mpfr_ptr, mpfr_srcptr, mpfr_srcptr, mpfr_rnd_t);
43 extern int mpfr_sub(mpfr_ptr, mpfr_srcptr, mpfr_srcptr, mpfr_rnd_t);
44 extern int mpfr_cmp_ui(mpfr_srcptr, unsigned long);
45 extern double mpfr_get_d(mpfr_srcptr, mpfr_rnd_t);
46 extern const char *mpfr_get_version(void);
47 #endif
48
49 static inline int is_power_of_two_u64(uint64_t x) {
50     return x && ((x & (x - 1)) == 0);
51 }
52
53 typedef struct {
54     uint64_t checked;
55     uint64_t failures;
56     double min_lower;
57     double min_upper;
58     uint64_t min_lower_m, min_lower_n;
59     uint64_t min_upper_m, min_upper_n;
60     uint64_t first_fail_m, first_fail_guess;
61 } Stats;
62
63 int main(int argc, char **argv) {
64     uint64_t start = 1;
65     uint64_t limit = (1ULL << 20);
66     int bits = 192;
67     if (argc > 1) start = strtoull(argv[1], NULL, 0);
68     if (argc > 2) limit = strtoull(argv[2], NULL, 0);
69     if (argc > 3) bits = atoi(argv[3]);
70     if (start < 1 || limit <= start) {
71         fprintf(stderr, "bad range\n");
72         return 2;
73     }
74
75     int nt = omp_get_max_threads();
76     Stats *stats = calloc((size_t)nt, sizeof(Stats));
77     if (!stats) return 3;
78     for (int i = 0; i < nt; ++i) {
79         stats[i].min_lower = INFINITY;
80         stats[i].min_upper = INFINITY;
81     }
82
83     const long double theta = logl(3.0L) / logl(2.0L);
84     double t0 = omp_get_wtime();
85
86     mpfr_t two, three, ln2_lo, ln2_hi, ln3_lo, ln3_hi;
87     mpfr_init2(two, bits); mpfr_init2(three, bits);
88     mpfr_init2(ln2_lo, bits); mpfr_init2(ln2_hi, bits);

```



```

89  mpfr_init2(ln3_lo, bits); mpfr_init2(ln3_hi, bits);
90  mpfr_set_ui(two, 2, MPFR_RNDN);
91  mpfr_set_ui(three, 3, MPFR_RNDN);
92  mpfr_log(ln2_lo, two, MPFR_RNDD);
93  mpfr_log(ln2_hi, two, MPFR_RNDU);
94  mpfr_log(ln3_lo, three, MPFR_RNDD);
95  mpfr_log(ln3_hi, three, MPFR_RNDU);
96
97  #pragma omp parallel
98  {
99      int tid = omp_get_thread_num();
100      Stats *st = &stats[tid];
101      mpfr_t xm, xn, xnp1;
102      mpfr_t logm_lo, logm_hi, logn_hi, lognp1_lo;
103      mpfr_t lhs_lo, lhs_hi, rhs_lo, rhs_hi, dl, du;
104      mpfr_init2(xm, bits); mpfr_init2(xn, bits); mpfr_init2(xnp1, bits);
105      mpfr_init2(logm_lo, bits); mpfr_init2(logm_hi, bits);
106      mpfr_init2(logn_hi, bits); mpfr_init2(lognp1_lo, bits);
107      mpfr_init2(lhs_lo, bits); mpfr_init2(lhs_hi, bits);
108      mpfr_init2(rhs_lo, bits); mpfr_init2(rhs_hi, bits);
109      mpfr_init2(dl, bits); mpfr_init2(du, bits);
110
111  #pragma omp for schedule(static)
112  for (uint64_t m = start; m < limit; ++m) {
113      if (is_power_of_two_u64(m)) continue;
114
115      /* This is only a locator. The MPFR inequalities below are the proof. */
116      uint64_t guess = (uint64_t)floorl(expl(theta * logl((long double)m)));
117
118      mpfr_set_ui(xm, (unsigned long)m, MPFR_RNDN);
119      mpfr_log(logm_lo, xm, MPFR_RNDD);
120      mpfr_log(logm_hi, xm, MPFR_RNDU);
121      mpfr_mul(lhs_lo, logm_lo, ln3_lo, MPFR_RNDD);
122      mpfr_mul(lhs_hi, logm_hi, ln3_hi, MPFR_RNDU);
123
124      int ok = 0;
125      uint64_t certn = 0;
126      double dl_d = 0.0, du_d = 0.0;
127      static const int offsets[7] = {0, -1, 1, -2, 2, -3, 3};
128      for (int oi = 0; oi < 7; ++oi) {
129          int off = offsets[oi];
130          if (off < 0 && guess < (uint64_t)(-off) + 1) continue;
131          uint64_t n = (uint64_t)((int64_t)guess + off);
132          if (n < 1) continue;
133
134          mpfr_set_ui(xn, (unsigned long)n, MPFR_RNDN);
135          mpfr_set_ui(xnp1, (unsigned long)(n + 1), MPFR_RNDN);
136
137          /* Directed lower bound for log(m)log(3)-log(n)log(2). */
138          mpfr_log(logn_hi, xn, MPFR_RNDU);
139          mpfr_mul(rhs_hi, logn_hi, ln2_hi, MPFR_RNDU);
140          mpfr_sub(dl, lhs_lo, rhs_hi, MPFR_RNDD);
141          if (mpfr_cmp_ui(dl, 0) <= 0) continue;
142
143          /* Directed lower bound for log(n+1)log(2)-log(m)log(3). */
144          mpfr_log(lognp1_lo, xnp1, MPFR_RNDD);
145          mpfr_mul(rhs_lo, lognp1_lo, ln2_lo, MPFR_RNDD);
146          mpfr_sub(du, rhs_lo, lhs_hi, MPFR_RNDD);
147          if (mpfr_cmp_ui(du, 0) <= 0) continue;

```

```

148
149     ok = 1;
150     certn = n;
151     dl_d = mpfr_get_d(dl, MPFR_RNDD);
152     du_d = mpfr_get_d(du, MPFR_RNDD);
153     break;
154 }
155
156 if (!ok) {
157     ++st->failures;
158     if (!st->first_fail_m || m < st->first_fail_m) {
159         st->first_fail_m = m;
160         st->first_fail_guess = guess;
161     }
162     continue;
163 }
164
165 ++st->checked;
166 if (dl_d < st->min_lower) {
167     st->min_lower = dl_d;
168     st->min_lower_m = m;
169     st->min_lower_n = certn;
170 }
171 if (du_d < st->min_upper) {
172     st->min_upper = du_d;
173     st->min_upper_m = m;
174     st->min_upper_n = certn;
175 }
176 }
177
178 mpfr_clear(xm); mpfr_clear(xn); mpfr_clear(xnp1);
179 mpfr_clear(logm_lo); mpfr_clear(logm_hi);
180 mpfr_clear(logn_hi); mpfr_clear(lognp1_lo);
181 mpfr_clear(lhs_lo); mpfr_clear(lhs_hi);
182 mpfr_clear(rhs_lo); mpfr_clear(rhs_hi);
183 mpfr_clear(dl); mpfr_clear(du);
184 }
185
186 Stats total = {0};
187 total.min_lower = INFINITY;
188 total.min_upper = INFINITY;
189 for (int i = 0; i < nt; ++i) {
190     Stats *s = &stats[i];
191     total.checked += s->checked;
192     total.failures += s->failures;
193     if (s->min_lower < total.min_lower) {
194         total.min_lower = s->min_lower;
195         total.min_lower_m = s->min_lower_m;
196         total.min_lower_n = s->min_lower_n;
197     }
198     if (s->min_upper < total.min_upper) {
199         total.min_upper = s->min_upper;
200         total.min_upper_m = s->min_upper_m;
201         total.min_upper_n = s->min_upper_n;
202     }
203     if (s->first_fail_m &&
204         (!total.first_fail_m || s->first_fail_m < total.first_fail_m)) {
205         total.first_fail_m = s->first_fail_m;
206         total.first_fail_guess = s->first_fail_guess;

```

```

207     }
208 }
209
210 double elapsed = omp_get_wtime() - t0;
211 printf("mpfr_version=%s\n", mpfr_get_version());
212 printf("range=[%" PRIu64 ",%" PRIu64 "]\n", start, limit);
213 printf("precision_bits=%d\nthreads=%d\n", bits, nt);
214 printf("checked_nonpowers=%" PRIu64 "\n", total.checked);
215 printf("failures=%" PRIu64 "\n", total.failures);
216 printf("min_lower_log_margin=%.17g at (m,n)=(%" PRIu64 ",%" PRIu64 ")\n",
217        total.min_lower, total.min_lower_m, total.min_lower_n);
218 printf("min_upper_log_margin=%.17g at (m,n)=(%" PRIu64 ",%" PRIu64 ")\n",
219        total.min_upper, total.min_upper_m, total.min_upper_n);
220 if (total.failures) {
221     printf("first_failure=(m,guess)=(%" PRIu64 ",%" PRIu64 ")\n",
222            total.first_fail_m, total.first_fail_guess);
223 }
224 printf("elapsed_seconds=%.6f\n", elapsed);
225
226 mpfr_clear(two); mpfr_clear(three);
227 mpfr_clear(ln2_lo); mpfr_clear(ln2_hi);
228 mpfr_clear(ln3_lo); mpfr_clear(ln3_hi);
229 free(stats);
230 return total.failures ? 1 : 0;
231 }

```

Listing 20: Directed-rounding verifier (C99, MPFR, OpenMP).

I.4 Build and invocation

On a system with the MPFR development headers installed, the whole build and a single full-range run are:

```

gcc -O3 -fopenmp mpfr_scan.c -lmpfr -lm -o mpfr_scan
OMP_NUM_THREADS=5 ./mpfr_scan 1 67108864 192

```

Listing 21: Build and invocation (POSIX shell).

The recorded audit instead used sixteen shorter invocations over half-open subranges, so that each chunk produces an independently inspectable log. The exit status is nonzero exactly when some input failed to be certified, which makes the chunked run scriptable without parsing the output.

J Chunk audit and reproducibility

Status: [\[software computation\]](#) — the numbers in this appendix are the output of the verifier of Appendix I. They are reproducible data at a software trust level, not kernel certificates; the kernel-verified exponent bound remains $x \geq 12$.

There are $2^{26} - 1 = 67,108,863$ positive integers below 2^{26} , of which exactly 26 are powers of two, namely 2^0 through 2^{25} . The verifier therefore tests $67,108,863 - 26 = 67,108,837$ inputs. The range

was split at multiples of 2^{22} into one initial chunk $[1, 2^{22})$ and fifteen chunks of width 2^{22} . The canonical chunk log holds the full standard output of every invocation and has SHA-256 digest

dcaa7f3b7e5be44acaee9015a2045da14ee42d8fa2d96b9619cdd44c072a6fd7

Table 16 reports, for each chunk, the exact scanned integer interval, the number of non-powers certified, and the smallest directed lower and upper logarithmic margins recorded in that chunk. Every chunk returned a failure count of zero.

Table 16: Directed-rounding audit through 2^{26} .

Chunk	Integer range	Checked	Smallest lower gap	Smallest upper gap
0	1–4,194,303	4,194,281	1.1539×10^{-17}	3.0076×10^{-17}
1	4,194,304–8,388,607	4,194,303	4.6319×10^{-20}	1.3515×10^{-19}
2	8,388,608–12,582,911	4,194,303	1.2811×10^{-18}	1.3515×10^{-19}
3	12,582,912–16,777,215	4,194,304	4.6319×10^{-20}	2.9082×10^{-19}
4	16,777,216–20,971,519	4,194,303	1.1898×10^{-18}	1.3515×10^{-19}
5	20,971,520–25,165,823	4,194,304	5.4212×10^{-19}	1.5111×10^{-19}
6	25,165,824–29,360,127	4,194,304	4.6319×10^{-20}	1.7669×10^{-19}
7	29,360,128–33,554,431	4,194,304	1.4439×10^{-19}	5.4132×10^{-21}
8	33,554,432–37,748,735	4,194,303	4.2839×10^{-19}	1.3515×10^{-19}
9	37,748,736–41,943,039	4,194,304	3.9292×10^{-19}	1.0460×10^{-19}
10	41,943,040–46,137,343	4,194,304	1.9177×10^{-19}	1.5111×10^{-19}
11	46,137,344–50,331,647	4,194,304	1.0117×10^{-19}	3.4145×10^{-19}
12	50,331,648–54,525,951	4,194,304	5.6927×10^{-20}	1.2322×10^{-19}
13	54,525,952–58,720,255	4,194,304	3.6447×10^{-20}	1.1345×10^{-19}
14	58,720,256–62,914,559	4,194,304	6.7779×10^{-20}	5.4132×10^{-21}
15	62,914,560–67,108,863	4,194,304	5.0923×10^{-20}	3.9197×10^{-20}
total		67,108,837		

J.1 Extremal records

The global lower-gap record is the chunk 13 value 3.6447×10^{-20} , attained at

$$(m, n) = (58,570,438, 2,048,703,434,093),$$

where a separate 100-digit evaluation gives $m^\vartheta - n = 1.07725158750919017548 \dots \times 10^{-7}$. The global upper-gap record is the chunk 7 value 5.4132×10^{-21} , attained at

$$(m, n) = (29,411,704, 687,581,893,443),$$

where $n + 1 - m^\vartheta = 5.36974926710988874953 \dots \times 10^{-9}$. The same upper margin reappears in chunk 14 at the scaled pair $(58,823,408, 2,062,745,680,331)$; this is not a coincidence but the identity $(2m)^\vartheta = 3m^\vartheta$, which multiplies the input by two and its ϑ -power by three while leaving the logarithmic margin unchanged.

Both extremal inequalities were reevaluated independently at 100 decimal digits. That secondary check is not an independent proof system — it shares the same floating-point philosophy — but it does guard against transcription and aggregation mistakes, and the certified margins sit many orders of magnitude above the nominal 192-bit rounding scale of the logarithmic operations.

The 2^{20} run of Appendix F and the present run are independent implementations, one using interval arithmetic in `mpmath` and the other MPFR directed rounding, and they agree on the overlapping range $1 \leq m < 2^{20}$: no failures. Their recorded extremal metadata are not directly comparable, because the chunk boundaries differ.

J.2 Reproducibility checklist

To reproduce the finite claim independently:

1. obtain MPFR 4.2.x and GMP on a 64-bit platform;
2. save the source of Appendix I and verify its SHA-256 digest;
3. compile with optimization and OpenMP, or delete the OpenMP pragmas for a serial run;
4. run the sixteen half-open intervals of Table 16 at 192 bits;
5. verify that every invocation reports `failures=0`;
6. sum the reported `checked_nonpowers` and obtain 67,108,837;
7. compare the two global extremal records above;
8. optionally rerun at 256 or 320 bits, and on a second MPFR build;
9. for a theorem-grade result, translate the intervals into exact rational certificates and replay them in Lean rather than trusting the executable.

The finite statement is deterministic. Thread count changes execution time and the reduction order that produces the recorded minimum metadata, but it cannot change whether an individual directed inequality is certified. Step nine is the one that would change the status of the result rather than its range, and it is the reason the scan is presented here as data for certificate generation — in the sense of Section 27 — and not as an extension of the kernel-verified zero-free region.

K Classical transcendence inputs and their exact roles

Status: `[classical/open background]` — this appendix states external theorems, not new ones. Exactly one of them, Gelfond–Schneider, is formalized in this repository (Theorem 4.11); the remaining ones enter only through paper proofs, which are tagged `[paper proof in this report]` where they appear. The corpus does not introduce any of them as an axiom: where a conjectural or unformalized statement is needed in Lean it is defined as a proposition and the implication is proved conditionally on it. Deleting six exponentials, five exponentials, Baker’s theorems, or Roy’s theorem from the report would change which paper results stand, and would change none of the kernel-verified results catalogued in Appendix A.

The purpose of this appendix is auditability. A reader who wants to know how much of the report is elementary consequence of the repository hypotheses, and how much rests on major transcendence theory, should be able to answer that question by looking in one place. For each classical input below we give the statement in the form actually used, the exact places where the report uses it,

and the exact reason it stops short of Conjecture 2.1. The last reason is the same in every case, seen from a different angle: every input listed here constrains a *linear* configuration of logarithms or an *algebraic* exponent, while the hypothetical counterexample is a *bilinear* relation between logarithms with a transcendental exponent.

Gelfond–Schneider (Hilbert’s seventh problem)

Statement. Let $\alpha, \beta \in \overline{\mathbb{Q}}$ with $\alpha \neq 0, 1$ and $\beta \notin \mathbb{Q}$. Then every value of α^β is transcendental [4, 9, 3]. [classical/open background]

Role here. This is the one classical transcendence theorem the corpus proves rather than cites; the Lean proof follows the classical exposition [4] and is recorded as Theorem 4.11. It is used twice. First, a nonintegral solution x is shown irrational by the elementary rational-exponent argument, and then must be transcendental because 2^x is an algebraic number; this is the algebraic–integral dichotomy of Corollary 4.12, and it is what makes the conditional primitive generator β and the odd-core exponent $\log_2 w$ transcendental throughout Section 5 and Section 6. Second, $\vartheta = \log_2 3$ is irrational, hence transcendental, because $2^\vartheta = 3$.

Why it does not settle Conjecture 2.1. The theorem speaks only about algebraic irrational exponents. Once the dichotomy has removed algebraic exponents, the remaining hypothetical exponent is transcendental, and Gelfond–Schneider has nothing further to say about 2^x for transcendental x . The theorem closes the easy half of the problem and defines the hard half.

Six exponentials

Statement. Let $u_1, u_2, u_3 \in \mathbb{C}$ be linearly independent over $\mathbb{Q}$ and let $v_1, v_2 \in \mathbb{C}$ be linearly independent over $\mathbb{Q}$. Then at least one of the six numbers

$$e^{u_i v_j} \quad (1 \leq i \leq 3, 1 \leq j \leq 2)$$

is transcendental [10, 11, 12]. [classical/open background]

Role here. It is the analytic engine behind every exact spectrum statement in the report. The saturated algebraic-base spectrum — for multiplicatively independent $a, b \in \mathbb{Q}_{>0}$ with $a^x, b^x \in \mathbb{Q}$ and x irrational, $\alpha^x \in \overline{\mathbb{Q}}$ if and only if $\alpha \in \Gamma(a, b)$ — is the case with rows $\log a, \log b, \log \alpha$ and columns $1, x$. The conditional classification of the symmetric constant $E_3 = 3^\vartheta$ is the case with rows $1, x, \vartheta$ and columns $\log 2, \log 3$, and the classification of $E_2 = 2^\eta$ is its mirror image under $\vartheta \mapsto \eta = 1/\vartheta$. The ordinary form of the unconditional E_3/E_2 dichotomy uses the same configuration. The kernel-verified 3-smooth classification of auxiliary bases and the three-base theorem $2^x, 3^x, 5^x \in \mathbb{Z} \Rightarrow x \in \mathbb{Z}$ are the formalized integer-level shadows of this theorem; the Lean proofs proceed by interpolation determinants and do not cite it.

Why it does not settle Conjecture 2.1. Six exponentials needs three independent directions on one side of the configuration. The two-base problem supplies exactly two logarithmic directions, $\log 2$ and $\log 3$, and every auxiliary base built from 2 and 3 has logarithm in their rational span, so it creates no third column. The remaining 2×2 configuration of Section 7 is precisely the four exponentials conjecture, which is open.

Five exponentials

Statement. Let u_1, u_2 be linearly independent over $\mathbb{Q}$, let v_1, v_2 be linearly independent over $\mathbb{Q}$, and let $\gamma \in \overline{\mathbb{Q}}$ with $\gamma \neq 0$. Then at least one of the five numbers

$$e^{u_1 v_1}, \quad e^{u_1 v_2}, \quad e^{u_2 v_1}, \quad e^{u_2 v_2}, \quad e^{\gamma u_2 / u_1}$$

is transcendental [11, 8]. [classical/open background]

Role here. It supplies the two forbidden exponential rays of the hypothetical world: for a nonintegral solution x and any nonzero algebraic γ , both $e^{\gamma x}$ and e^{γ^ϑ} are transcendental. For the first, apply the theorem to the independent pairs $(1, x)$ and $(\log 2, \log 3)$, whose four rectangular exponentials are the algebraic numbers $2, 3, M, A$, leaving the fifth transcendental; the second is the same argument with the roles of the two logarithms exchanged. Formalizing it would also upgrade the denominator-controlled auxiliary-base theorem in `RationalThirdBase` from a controlled denominator to an unrestricted rational output, which is why it appears as a research priority in Section 27.

Why it does not settle Conjecture 2.1. The theorem does not remove the fifth exponential, so it does not prove four exponentials. More concretely, the natural auxiliary quantities in this problem are $p^x = e^{x \log p}$ for a prime p , whose coefficient $\log p$ is transcendental rather than an admissible algebraic γ . The forbidden rays therefore never convert a prime divisor of the integer outputs into a contradiction, as Section 26 records.

Baker’s theorem on linear forms in logarithms

Statement. Let $\gamma_1, \dots, \gamma_n$ be nonzero algebraic numbers of degree at most d and height at most H , and let $\alpha_1, \dots, \alpha_n$ be algebraic numbers of height at most B , not all zero. If

$$\Lambda = \alpha_1 \log \gamma_1 + \dots + \alpha_n \log \gamma_n \neq 0,$$

then $|\Lambda| > B^{-C}$ for an effectively computable $C > 0$ depending only on n, d, H and the degree of the α_i [3, 10]. [classical/open background]

Role here. The effective form is used as a measuring stick rather than as a step in any proof. In the rational-approximation analysis of Section 26 it fixes the scale of what is available: Baker-type lower bounds are polynomial in the relevant heights, whereas the integer gaps produced by the hypotheses only assert $|\varepsilon| \gtrsim 2^{-p}$, exponentially weaker, so the two scales are consistent and no contradiction arises. It is also the theorem that effective irrationality measures for ϑ would refine.

Why it does not settle Conjecture 2.1. Baker’s theory bounds forms whose coefficients are algebraic. The decisive relation under failure is $\log M \log 3 = \log A \log 2$, and its radical form is $d \log w \log 3 = (\log A - a \log 3) \log 2$: each side is a product of two logarithms, so the coefficient of every logarithm is itself a logarithm. Expanded over prime factorizations $m = \prod_p p^{m_p}$ and $A = \prod_p p^{a_p}$, failure is the nontrivial vanishing of

$$D = \sum_p (m_p \log 3 - a_p \log 2) \log p,$$

an integer-coefficient linear form in the products $\log p \cdot \log 2$ and $\log p \cdot \log 3$. No lower-bound technology exists for such forms, and even the $\mathbb{Q}$ -linear independence of $\{\log p \cdot \log q\}$ is open; Section 12 turns exactly that independence into a sufficient condition.

Baker’s homogeneous theorem

Statement. Let $\gamma_1, \dots, \gamma_n$ be nonzero algebraic numbers. If $\log \gamma_1, \dots, \log \gamma_n$ are linearly independent over $\mathbb{Q}$, then $1, \log \gamma_1, \dots, \log \gamma_n$ are linearly independent over $\mathbb{Q}$. The homogeneous consequence used here is the shorter one: logarithms of algebraic numbers that are linearly independent over $\mathbb{Q}$ are linearly independent over $\overline{\mathbb{Q}}$ [3, 10]. [classical/open background]

Role here. It is the bridge from a rational-independence hypothesis to the $\overline{\mathbb{Q}}$ -independence hypothesis that Roy’s theorem requires. In the strong spectrum, the three logarithms $\log a, \log b, \log \alpha$

are first shown $\mathbb{Q}$ -independent by an elementary saturation argument, and Baker’s homogeneous theorem upgrades them to $\overline{\mathbb{Q}}$ -independence; the same upgrade produces the logarithmic independence used in the strong forms of the E_3 and E_2 classifications. The ordinary spectrum does not need it, since rational independence already suffices for six exponentials.

Why it does not settle Conjecture 2.1. It is a statement about linear relations among logarithms. No independence statement about the family $\{\log p\}$, over $\mathbb{Q}$ or over $\overline{\mathbb{Q}}$, implies anything about the family of products $\{\log p \cdot \log q\}$, and it is the products that carry the obstruction.

Roy’s strong and sharp six exponentials theorems

Statement. Let $\mathcal{L}$ be the $\overline{\mathbb{Q}}$ -vector space spanned by 1 together with all logarithms of nonzero algebraic numbers. If u_1, u_2, u_3 are linearly independent over $\overline{\mathbb{Q}}$ and v_1, v_2 are linearly independent over $\mathbb{Q}$, then at least one of the six products $u_i v_j$ lies outside $\mathcal{L}$. The sharp form replaces the independence hypotheses by a rank condition on a matrix whose entries are linear forms in logarithms of algebraic numbers [8, 11]. [classical/open background]

Role here. It produces the strengthened conclusions only. Concretely: the strong auxiliary-base spectrum, which asserts $x \log \alpha \notin \mathcal{L}$ for $\alpha \notin \Gamma(a, b)$ rather than merely $\alpha^x \notin \overline{\mathbb{Q}}$; the strong form of the symmetric-constant dichotomy, which places $\vartheta \log 3$ and $\eta \log 2$ outside $\mathcal{L}$; and the log-strong transcendence clauses in the second-iterate kernel and depth-three tower theorems.

Why it does not settle Conjecture 2.1. Every ordinary transcendence conclusion in the report survives the removal of this theorem; only the “lies outside $\mathcal{L}$ ” clauses disappear. Structurally it inherits the six exponentials dimension count and still demands three independent directions on one side, so it does not reach the 2×2 configuration where the conjecture actually lives.

Lindemann–Weierstrass

Statement. If $\alpha_1, \dots, \alpha_n$ are distinct algebraic numbers, then $e^{\alpha_1}, \dots, e^{\alpha_n}$ are linearly independent over $\overline{\mathbb{Q}}$. The special case used here is Hermite–Lindemann: for algebraic $\alpha \neq 0$ the value e^α is transcendental, equivalently every nonzero value of $\log \gamma$ is transcendental for algebraic $\gamma \neq 0, 1$ [3]. [classical/open background]

Role here. It is the standing background fact that each $\log p$, and in particular $\log 2$ and $\log 3$, is transcendental. That fact is what makes the coefficients in the log-product form D above transcendental rather than algebraic, and therefore what places the wall outside Baker’s reach; it is also the precise content of the audit warning against treating $\log 2$ and $\log 3$ as algebraic coefficients (Remark L.3).

Why it does not settle Conjecture 2.1. Lindemann–Weierstrass constrains exponentials of *algebraic* arguments. Under failure the exponent x is transcendental and the arguments $x \log 2$ and $x \log 3$ are transcendental, so none of the four corner exponentials $2, 3, 2^x, 3^x$ is within its scope, and it gives no information whatever about products of two logarithms.

K.1 Logical dependency graph

The table records, for each family of results in the report, the classical input it actually consumes. The rows whose input column reports no classical theorem are elementary or kernel-verified, and depend on no transcendence theory beyond the Gelfond–Schneider dichotomy that the corpus already formalizes.

Table 17: Classical transcendence input consumed by each family of results.

Result	Required classical input
Nonintegral solution is transcendental	Rational-exponent lemma and Gelfond–Schneider; both kernel-verified
Reduction to real four exponentials	None beyond $\mathbb{Q}$ -independence of $1, x$ and of $\log 2, \log 3$; kernel-verified
3-smooth auxiliary-base classification and the three-base theorem	None; interpolation determinants, kernel-verified
Exact algebraic-base spectrum for an independent pair	Six exponentials
Strong $\mathcal{L}$ spectrum	Exact row independence, Baker’s homogeneous theorem, and Roy
Valuation lattice and Smith normal form of rational outputs	Exact spectrum and unique factorization
Second-iterate kernel: dimension and the four branch types	Exact spectrum and elementary prime valuations
Kernel invariance across all nonintegral solutions	Exact spectrum applied twice
Möbius orbit criterion	Kernel definition and real logarithms
Depth-three tower theorem	Exact spectrum and kernel dimension; Roy only for the strong form
Base-2 tower equivalence	Natural-base algebraicity classification and smooth branch algebra
Symmetric constants E_3 and E_2	Six exponentials; Baker’s homogeneous theorem and Roy for the strong form
Two forbidden exponential rays $e^{\gamma x}, e^{\gamma \vartheta}$	Five exponentials
Transcendence of the coefficients $\log p$ in the log-product form	Lindemann–Weierstrass, in its Hermite–Lindemann case

L Common false shortcuts and audit warnings

Status: [classical/open background] — this appendix records negative facts and standard cautions. Nothing in it is a new result, and nothing in it changes the status of any statement elsewhere in the report.

The problem attracts several plausible but invalid one-line arguments. Recording them is part of the audit trail: each is a specific way in which a proposed proof can be circular, can silently assume the conclusion, or can mistake numerical evidence for exclusion. Future paper and formal work should be checked against this list before a short argument is believed.

Remark L.1 (Fractional-part error). Writing $x = n + \alpha$ with $n \in \mathbb{Z}$ and $0 < \alpha < 1$, and assuming $2^x \in \mathbb{Z}$, does *not* imply $2^\alpha \in \mathbb{Z}$. It implies only $2^\alpha \in \mathbb{Z}[1/2]$, and an irrational α is not an ordinary radical exponent $1/q$. The integrality of the output does not descend to the fractional part.

Remark L.2 (Rational-root circularity). From $M^\vartheta \in \mathbb{Z}$ one cannot conclude that ϑ is rational, nor that M is a perfect power. Gelfond–Schneider constrains algebraic irrational exponents, whereas ϑ is transcendental. An argument that quietly treats ϑ as a rational exponent has assumed the conclusion.

Remark L.3 (Baker is linear, not bilinear). The relation $\log M \log 3 = \log A \log 2$ is not a Baker linear form. Treating $\log 2$ and $\log 3$ as algebraic coefficients is invalid: by Hermite–Lindemann both are transcendental. The same objection applies to the radical form $d \log w \log 3 = (\log A - a \log 3) \log 2$ and to the prime-expanded log-product form of Section 12.

Remark L.4 (No modular real exponent). The expression $2^x \bmod p$ has no canonical meaning for a general real x . A p -adic exponential requires a p -adic exponent and a unit in its convergence domain; the real equality $2^x \in \mathbb{Z}$ supplies neither. Congruence arguments must therefore be stated for the integer outputs M and A , never for the real exponent.

Remark L.5 (Candidate quotients need not be candidates). If $m_1, m_2 \in \mathcal{C}$ and $m_1 \mid m_2$, it does not follow that $m_2/m_1 \in \mathcal{C}$, because the corresponding quotient of the outputs need not be integral. Multiplication is safe; division requires a separate valuation check. This is the reason $\mathcal{C}$ is treated as a multiplicative monoid rather than as a divisor-closed set.

Remark L.6 (Density alone is compatible). The orbit $\{kx\}$ is dense modulo one for irrational x , but the associated rational points have exponential height. A sequence of $O(\log H)$ points up to height H is compatible with generic transcendental-curve counting, so density by itself yields no contradiction without a denominator-sensitive refinement.

Remark L.7 (Numerical tolerance is not exclusion). Computing M^ϑ in floating point and comparing its distance to the nearest integer against a fixed epsilon cannot certify a zero-free region, however wide. Outward directed interval bounds or an exact certificate are required, which is why the finite scans in this report are tagged [\[software computation\]](#) and the kernel-verified exponent bound remains $x \geq 12$.

Remark L.8 (Algebraicity of a power is branch-sensitive). From $z^d \in \overline{\mathbb{Q}}$ with an integer $d > 0$ one may conclude $z \in \overline{\mathbb{Q}}$, but from $z^x \in \overline{\mathbb{Q}}$ with transcendental x no such algebraic closure argument is available. The finite-degree argument uses the polynomial $X^d - z^d$, which has no analogue when the exponent is not an integer.

Every warning above was re-derived independently during the preparation of the present unified report: each was proposed at some point as a short route to the conjecture, and each failed on exactly the ground stated here.

M Logarithmic-energy convergence for projected lattice triangles

We record the technical point needed in the triangular determinant limit. Let $\lambda_1, \lambda_2 > 0$ and assume their ratio has finite irrationality type: there are $c, \tau > 0$ such that

$$|r\lambda_1 + s\lambda_2| \geq c(1 + |r| + |s|)^{-\tau}$$

for all nonzero $(r, s) \in \mathbb{Z}^2$. Let $t_{1,N} < \cdots < t_{N,N}$ be the first N values of $a\lambda_1 + b\lambda_2$ with $a, b \in \mathbb{N}_0$, and set $\tilde{t}_{i,N} = t_{i,N}/\sqrt{N}$. Then the empirical measures converge to

$$d\mu(t) = \frac{2t}{U^2} dt, \quad U = \sqrt{2\lambda_1\lambda_2},$$

and

$$\frac{2}{N^2} \sum_{i < j} \log |\tilde{t}_{j,N} - \tilde{t}_{i,N}| \longrightarrow \iint \log |x - y| d\mu(x) d\mu(y). \quad (485)$$

The weak convergence follows from a Riemann-sum lattice count in the triangle $\lambda_1 a + \lambda_2 b \leq U\sqrt{N}$. For the singular kernel, truncate $\log |x - y|$ below at $-M$. Weak convergence handles the truncated kernel. Uniform integrability of the negative tail follows by counting difference vectors (r, s) in thin strips around $r\lambda_1 + s\lambda_2 = 0$; the finite-type lower bound prevents an individual difference from being smaller than a fixed power of N , while the strip count shows that the proportion below ε is $O(\varepsilon)$ up to a power-saving discrepancy term. Integrating the distribution function of $-\log |x - y|$ and then sending $M \rightarrow \infty$ proves (485).

For $(\lambda_1, \lambda_2) = (\log 2, \log 3)$, the required finite type follows from linear forms in logarithms. For $(1, \beta)$, it follows from $\beta = \log A / \log 2$ and the multiplicative independence of A and 2.

N Integral evaluations used in the determinant constants

N.1 The beta-two logarithmic energy

For the probability density $f(x) = 2x$ on $[0, 1]$,

$$\begin{aligned} \iint_0^1 4xy \log |x - y| dx dy &= 8 \int_0^1 \int_0^x xy \log(x - y) dy dx \\ &= 8 \int_0^1 x^3 \left(\frac{1}{2} \log x - \frac{3}{4} \right) dx \\ &= -\frac{7}{4}. \end{aligned}$$

Scaling by U adds $\log U$, proving Equation (228).

N.2 The simplex assignment constant

A coordinate of a point uniformly distributed on the standard two-simplex has quantile $q(s) = 1 - \sqrt{1 - s}$. Therefore

$$\begin{aligned} J &= \int_0^1 q(s)q(1 - s) ds \\ &= \int_0^1 \left(1 - \sqrt{s} - \sqrt{1 - s} + \sqrt{s(1 - s)} \right) ds \\ &= 1 - \frac{4}{3} + \frac{\pi}{8} = \frac{\pi}{8} - \frac{1}{3}. \end{aligned}$$

Its comonotone second-moment constant is $K = \int_0^1 q(s)^2 ds = 1/6$.

N.3 The balanced-rectangle energy

Let $Z = (X + Y)/2$ with independent uniform $X, Y \in [0, 1]$. The density of $S = X + Y$ is triangular, and the density of the difference of two independent copies of S is

$$h(t) = \begin{cases} \frac{2}{3} - t^2 + \frac{1}{2}t^3, & 0 \leq t \leq 1, \\ \frac{1}{6}(2 - t)^3, & 1 \leq t \leq 2. \end{cases}$$

A direct integration gives

$$\mathbb{E} \log |S - S'| = \frac{4}{3} \log 2 - \frac{25}{12}.$$

Since $Z - Z' = (S - S')/2$,

$$\mathbb{E} \log |Z - Z'| = \frac{1}{3} \log 2 - \frac{25}{12}.$$

Also $\mathbb{E} Z^2 = 7/24$.

O General two-prime version

The structural and determinant arguments extend verbatim to distinct primes p, q . If a nonintegral β satisfies $p^\beta, q^\beta \in \mathbb{Z}$, choose the least such β . Six Exponentials gives the same affine-rank collapse, and valuation minimality yields the exact monoid $\mathbb{N}_0 + \mathbb{N}_0\beta$. The triangular determinant then gives

$$g\left(\sqrt{\beta \log p \log q}\right) \geq 0.$$

This form may be useful for comparing the strength of the method across base pairs and for testing variational improvements.

P Finite determinant certificates

The asymptotic argument can be converted into finite certificates. For chosen finite node sets, a certificate consists of:

1. the exact integer or rational matrix entries;
2. sorted rational intervals for u_i, v_i and their differences;
3. an upper bound from Equation (223);
4. for each relevant prime, dual potentials for the assignment problem in Lemma 51.3, proving the claimed minimum;
5. a product-formula comparison showing the upper bound is smaller than the certified divisor.

Such certificates are compact enough for independent checking and eventual Lean verification. They also avoid trusting floating-point determinant computations, which are extremely ill-conditioned in the near-singular regime.

Q Source of the finite verifier

The following is the complete source used to establish Theorem 63.1. Its only non-rigorous arithmetic—the binary64 call used to suggest a starting integer—cannot affect the result: the subsequent directed-interval sign search is monotone and independently certifies the final pair of inequalities.

```
#!/usr/bin/env python3
"""Rigorous finite verification for the Alaoglu--Erdos problem.

For every integer A in [2, LIMIT], this proves that  $A^{(\log(3)/\log(2))}$ 
is an integer only when A is a power of two. All decisive comparisons use
Decimal directed rounding plus the positive atanh series for logarithms.
The ordinary binary64 approximation is used only to choose an initial nearby
integer; a monotone certified search repairs any inaccurate guess.
"""
from __future__ import annotations

from dataclasses import dataclass
from decimal import Decimal, Context, ROUND_FLOOR, ROUND_CEILING, localcontext
import json
import math
import sys
import time

LIMIT = int(sys.argv[1]) if len(sys.argv) > 1 else 100_000
PREC = 70
TERMS = 28

DOWN = Context(prec=PREC, rounding=ROUND_FLOOR)
UP = Context(prec=PREC, rounding=ROUND_CEILING)

@dataclass(frozen=True)
class Interval:
    lo: Decimal
    hi: Decimal

def _atanh_log_ratio(p: int, q: int, terms: int = TERMS) -> Interval:
    """Enclose  $2*\operatorname{atanh}(p/q)$ , for  $0 \leq p/q \leq 1/3$ ."""
    assert 0 <= p <= q // 3 + (1 if 3 * p == q else 0)
    if p == 0:
        z = Decimal(0)
        return Interval(z, z)

    with localcontext(DOWN):
        z_lo = Decimal(p) / Decimal(q)
        z2_lo = z_lo * z_lo
        power = z_lo
        total = Decimal(0)
        for j in range(terms):
            total = total + power / Decimal(2 * j + 1)
            power = power * z2_lo
        lo = Decimal(2) * total

    with localcontext(UP):
        z_hi = Decimal(p) / Decimal(q)
        z2_hi = z_hi * z_hi
        power = z_hi
        total = Decimal(0)
        for j in range(terms):
            total = total + power / Decimal(2 * j + 1)
            power = power * z2_hi
        partial = Decimal(2) * total
        # power is now  $z^{(2*terms+1)}$ . Since all omitted terms are positive,
        # replace every odd denominator by the first one and sum geometrically.
```

```

        rem = Decimal(2) * power / (
            Decimal(2 * terms + 1) * (Decimal(1) - z2_hi)
        )
        hi = partial + rem
    return Interval(lo, hi)

LOG2 = _atanh_log_ratio(1, 3)

def log_integer(n: int, terms: int = TERMS) -> Interval:
    """Rigorous interval for log(n), n >= 1, by binary range reduction."""
    assert n >= 1
    if n == 1:
        z = Decimal(0)
        return Interval(z, z)
    k = n.bit_length() - 1
    two_k = 1 << k
    #  $n/2^k$  in  $[1, 2)$ , and  $z=(y-1)/(y+1)$  lies in  $[0, 1/3)$ .
    local = _atanh_log_ratio(n - two_k, n + two_k, terms)
    with localcontext(DOWN):
        lo = Decimal(k) * LOG2.lo + local.lo
    with localcontext(UP):
        hi = Decimal(k) * LOG2.hi + local.hi
    return Interval(lo, hi)

LOG3 = log_integer(3)

def determinant_log_interval(log_a: Interval, b: int) -> Interval:
    """Enclose  $F(A, b) = \log(b) \log(2) - \log(A) \log(3)$ ."""
    log_b = log_integer(b)
    with localcontext(DOWN):
        lo = log_b.lo * LOG2.lo - log_a.hi * LOG3.hi
    with localcontext(UP):
        hi = log_b.hi * LOG2.hi - log_a.lo * LOG3.lo
    return Interval(lo, hi)

def certified_sign(log_a: Interval, b: int) -> tuple[int, Interval]:
    f = determinant_log_interval(log_a, b)
    if f.hi < 0:
        return -1, f
    if f.lo > 0:
        return +1, f
    # At this finite precision an interval meeting zero is not silently accepted.
    raise ArithmeticError(f"inconclusive interval at A-log={log_a}, b={b}: {f}")

def is_power_of_two(n: int) -> bool:
    return n > 0 and (n & (n - 1)) == 0

def main() -> None:
    start = time.monotonic()
    rho_approx = math.log(3.0) / math.log(2.0)
    min_abs_endpoint = Decimal('Infinity')
    max_repairs = 0
    checked = 0

    for a in range(2, LIMIT + 1):
        if is_power_of_two(a):
            continue
        checked += 1
        log_a = log_integer(a)
        guess = max(1, int(math.floor(math.exp(rho_approx * math.log(float(a))))))
        repairs = 0

```

```

# Find consecutive integers b,b+1 on opposite sides of the unique real root.
while True:
    s0, f0 = certified_sign(log_a, guess)
    s1, f1 = certified_sign(log_a, guess + 1)
    min_abs_endpoint = min(
        min_abs_endpoint,
        abs(f0.lo), abs(f0.hi), abs(f1.lo), abs(f1.hi)
    )
    if s0 < 0 < s1:
        break
    if s1 < 0:
        guess += 1
    elif s0 > 0:
        guess -= 1
        assert guess >= 1
    else:
        raise AssertionError((a, guess, s0, s1))
    repairs += 1
    if repairs > 100:
        raise AssertionError(f"unexpectedly bad initial approximation at A={a}")
max_repairs = max(max_repairs, repairs)

elapsed = time.monotonic() - start
result = {
    "limit": LIMIT,
    "checked_nonpowers": checked,
    "precision_decimal_digits": PREC,
    "atanh_terms": TERMS,
    "max_initial_guess_repairs": max_repairs,
    "minimum_absolute_interval_endpoint": str(min_abs_endpoint),
    "elapsed_seconds": elapsed,
    "conclusion": (
        f"For every integer 2 <= A <= {LIMIT}, "
        "A^(log(3)/log(2)) is integral only if A is a power of two."
    ),
}
print(json.dumps(result, indent=2, sort_keys=True))

if __name__ == '__main__':
    main()

```

R External literature reconnaissance

This appendix merges two literature-and-strategy surveys of the same corpus. It records what the surrounding literature does and does not supply, which imports are immediately usable, and which are techniques that would need adaptation. Two claims made in the source surveys are retracted here; both are flagged where they occur.

Citations are given in full in the text rather than through the bibliography, because the purpose of the appendix is provenance rather than argument: nothing below is used as a step in any proof in this report.

R.1 Effective measures for the two logarithms

The sharpest linear-independence measure for the relevant triple is Qiang Wu, *On the linear independence measure of logarithms of rational numbers*, Math. Comp. **72** (2003), 901–911, which

gives $\mu(1, \log 2, \log 3) \leq 7.6155$, improving Rhin’s 7.616. For the single logarithm, Wu and Wang, *J. Number Theory* **142** (2014), 264–273, give $\mu(\log 3) \leq 5.1163051$, improving Salikhov’s 5.125. For two-logarithm forms, Laurent–Mignotte–Nesterenko, *J. Number Theory* **55** (1995), 285–321, and Laurent, *Acta Arith.* **133** (2008), 325–348, give explicit $\log |b_1 \log 2 + b_2 \log 3| \geq -C(\log B)^2$ with C of the order of a few tens. On the non-archimedean side, Bugeaud–Laurent (1996) and Bugeaud (1999) are the two-logarithm workhorses, and K. C. Chim, *J. Number Theory* (2025), improves the $(\log B)^2$ dependence to $\log B$ for two p -adic logarithms.

Caution R.1 (Retracted claim). The first survey proposed that Wu’s measure would yield a *uniform* zero-free criterion replacing the per-value certificates of Section 25. That is wrong, and the error is instructive. Near-integrality of U^ϑ is not a rational approximation to ϑ ; it is the vanishing of the bilinear form $\log 3 \log U - \log 2 \log V$, which is exactly the wall of Section 7. No linear-in-logarithms measure touches it. The finite verification stays per-value.

What the measure does give unconditionally is an effective irrationality exponent $\mu_{\text{eff}}(\vartheta) \leq 8.62$, hence a polynomial recurrence $q_{n+1} \ll q_n^{7.62}$ for the convergent denominators of ϑ . That is the strongest known bound of its kind and supersedes the exponential estimate; the kernel-verified module `DenominatorGrowth` proves only the weaker exponential form, which remains the strongest [\[kernel-verified Lean\]](#) statement here.

R.2 Interpolation determinants with non-archimedean amplification

The closest existing realization of the amplification idea audited in Section 21 is Bombieri, *Effective Diophantine approximation on $\mathbb{G}_m$* , *Ann. Scuola Norm. Sup. Pisa* **20** (1993), 61–89, and Bombieri–Cohen, *ibid.* **24** (1997), 205–225, with the elementary account in *LMS Lecture Notes* **303** (2003), 41–62. These combine archimedean and non-archimedean interpolation determinants through the equivariant Thue–Siegel principle and are the only non-Baker route to effective results on the multiplicative group. Laurent’s interpolation-determinant papers (*Acta Arith.* **66** (1994), 181–199, and **133** (2008)) supply the determinant method in the form this report already uses.

This is a technique needing adaptation, not a plug-in. Its relevance is sharpened by Corollary 21.2: any amplification must obtain its archimedean bound from genuine cancellation, and Bombieri–Cohen is where that is done systematically.

R.3 Small-value estimates and the four-exponentials frontier

Roy, *An arithmetic criterion for the values of the exponential function*, *Acta Arith.* **97** (2001), 183–194, constructs a small-value criterion on $\mathbb{G}_a \times \mathbb{G}_m$ and proves it *equivalent* to Schanuel’s conjecture. His subsequent estimates — *Acta Arith.* **135** (2008), 357–393; *Int. J. Number Theory* **6** (2010), 919–956; *Mathematika* **59** (2013), 333–363 — and Nguyen–Roy, *Int. J. Number Theory* **12** (2016), 1273–1293, are the deepest partial progress toward the four-exponentials and Schanuel circle. Roy’s Strong Six Exponentials Theorem, *J. Number Theory* **41** (1992), 22–47, is the object whose formalization would most directly strengthen the algebraic-base spectrum of Section 13.

Butler, *Some cases of Wilkie’s conjecture*, *Bull. London Math. Soc.* **44** (2012), 642–660, shows that certain cases of Wilkie’s conjecture are equivalent to real forms of the three- and four-exponentials conjectures. That is a direct bridge between the counting route below and the transcendence core, and it was not previously cited in this corpus.

R.4 Counting rational points on the transcendental arc

A counterexample produces $\asymp T$ rational points

$$P_k = \left(\frac{M^k}{2^{r_k}}, \frac{A^k}{3^{r_k}} \right) = (2^{\{k\beta\}}, 3^{\{k\beta\}})$$

of logarithmic height at most T , lying simultaneously in a *fixed* rank-two multiplicative subgroup $\Gamma \subset \mathbb{G}_m^2(\overline{\mathbb{Q}})$ and on the *fixed* transcendental Pfaffian arc $\mathcal{C} = \{(u, u^\vartheta) : 1 \leq u < 2\}$. This unlikely-intersection formulation is the most promising framing the surveys contribute, and it is not used elsewhere in this report.

The counting technology is Binyamini–Novikov, *Ann. of Math.* **186** (2017), 237–275, and Binyamini–Novikov–Zak, *Ann. of Math.* **199** (2024), 795–821, which prove Wilkie’s conjecture for restricted elementary and Pfaffian structures with $(\log H)^\kappa$ counting; and, on the non-archimedean side, Binyamini–Cluckers–Novikov (*Duke Math. J.*), Cluckers–Comte–Loeser, *Forum Math. Pi* **3** (2015), and *Counting rational points on transcendental curves in valued fields*, arXiv:2506.19411 (2025). Boxall–Jones–Schmidt, *J. Eur. Math. Soc.* **24** (2022), 1567–1592, combine archimedean counting with p -adic methods.

The honest verdict is negative in a familiar way. Polylogarithmic counting gives $(\log H)^\kappa$ where a counterexample supplies $\asymp \log H$ points, so the framing reproduces the one-logarithm deficit of Section 22 in counting form: a third independent route arriving at the same shortfall. Moreover no cited theorem is sensitive to *which* primes divide the denominators, which is precisely what the compact-arc reformulation of Section 23 needs. The statement required — that a compact transcendental arc of nonzero curvature carries $o(H)$ points of height at most H whose first coordinate has 2-power denominator and second 3-power denominator — is not implied by any of them and would be a new theorem.

R.5 Multiplicative dependence, S -units, and rigidity

Bugeaud–Corvaja–Zannier, *Math. Z.* **243** (2003), 79–84, prove that for multiplicatively independent $a, b \geq 2$ and every $\varepsilon > 0$ one has $\log \gcd(a^n - 1, b^n - 1) \leq \varepsilon n$ for large n , and that this is sharp. Together with the Corvaja–Zannier subspace-theorem machinery and Evertse–Györy’s effective S -unit theory, this confirms that S -unit methods deliver finiteness, not nonexistence, and cannot supply the exponential-in- q divisibility the rational-approximation route would need.

For the $\times 2 \times 3$ side, Bourgain–Lindenstrauss–Michel–Venkatesh, *Ergodic Theory Dynam. Systems* **29** (2009), 1705–1722, give the only *effective* Furstenberg-type density result, with a density gap of order $(\log \log L)^{-\kappa}$. Shmerkin and Wu, *Ann. of Math.* **189** (2019), resolved Furstenberg’s intersection conjecture. None of these can force visits to targets shrinking like M^{-k} ; the target-width mismatch is fatal for a direct application, exactly as Section 23 anticipates. The effective Bourgain–Lindenstrauss–Michel–Venkatesh rate is the one usable quantitative handle and confirms, rather than defeats, the shrinking-target obstruction.

R.6 Kummer theory and the radical reformulation

Perucca–Sgobba–Tronto, *Int. J. Number Theory* **17** (2021), 1091–1110, and *Kummer theory for number fields via entanglement groups*, *Manuscripta Math.* **169** (2022), 251–270, compute degrees of

Kummer extensions exactly through entanglement groups, and Chan–Pajaziti–Perissinotto–Perucca, arXiv:2508.19211 (2025), complete Kneser’s theorem on additive relations among radicals. These make the algebraic half of the radical reformulation of Section 6 rigorous and formalizable. They do not see the special value ϑ : the mixed archimedean–radical invariant that the reformulation needs is supplied by no Kummer-theoretic result found.

R.7 Why the generic-power model theory excludes this exponent

Bays, Kirby and Wilkie, J. London Math. Soc. (2010), prove a Schanuel property for raising to an *exponentially transcendental* power: for such λ and multiplicatively independent $\bar{y}$ one has $\text{td}(\bar{y}, \bar{y}^\lambda) \geq n$. The result does not apply here, and the reason is worth recording precisely, because it explains a pattern rather than merely reporting a failure.

The exponent ϑ is exponentially *algebraic*. The pair $(z, w) = (\log 2, \vartheta)$ is a nonsingular solution of the system

$$e^z - 2 = 0, \quad e^{zw} - 3 = 0,$$

whose Jacobian determinant is $6 \log 2 \neq 0$. So ϑ lies in exactly the exceptional set that the generic-power theorems are stated to avoid. The same holds for every version of the problem in this report, since all of them involve the same exponent. Generic o-minimal and model-theoretic machinery should therefore be treated as a structural guide here, not as a black box expected to apply.

R.8 Prioritized directions, and what the ceiling changes

The second survey graded the located literature by how directly it attaches to the obstruction. Its ranking, compressed, was: highest priority to the rank-two multiplicative group combined with specialized o-minimal or algebraic covering and S -unit rigidity; to a Laurent-style interpolation-determinant redesign of the semigroup determinant of Section 20; and to tropical or p -adic determinant amplification with modern p -adic logarithmic-form control. High priority to the Corvaja–Zannier recurrence and power-sum machinery, once an integral quotient has been manufactured, and to the Mahler–Ridout–Corvaja–Zannier exponential-nearness criterion, which supplies a sharp conditional target: construct one non-Pisot algebraic power sequence exponentially close to the integers. Medium priority to the explicit logarithmic-form technology of Matveev, Laurent, Gouillon and Mignotte–Voutier, excellent for finite verification and auxiliary lemmas but unlikely to cross the bilinear barrier alone, and to integer-valued o-minimal function theory, methodologically close but not fitted to a sparse multiplicative grid. Low priority, as a direct attack, to Skolem–Mahler–Lech, generic-power model theory and general Zilber–Pink.

Every item in that survey was recorded as entirely absent from the corpus at the time, and the present report’s bibliography still does not carry most of them; they are listed here so that the gap is visible rather than implicit.

One entry of the ranking has since been settled negatively. The third highest-priority direction, tropical or p -adic determinant amplification, is retired at the term-wise level by Corollary 21.2: automatic divisibility can never contradict a correct archimedean bound, whatever the primes, nodes or columns. What survives of that direction is the cancellation problem localized by Proposition 21.3, which is a different and harder target than the one the ranking had in view.

R.9 Provenance of the conjecture

Alaoglu and Erdős proved that the ratio of consecutive superabundant numbers is prime, and, using the three-prime case that Siegel had claimed, that the ratio of consecutive colossally abundant numbers is a prime or a semiprime. Erdős and Nicolas later showed that exactly one, two, or four integers can attain the maximum in the colossally abundant construction, so a positive answer to the rational-power form of the conjecture bounds this to at most two and makes the ratio of consecutive colossally abundant numbers prime. Lagarias, *Amer. Math. Monthly* **109** (2002), 534–543, records the Robin/Riemann reformulation of the surrounding circle.

Caution R.2 (Unverified figure). The first survey quoted a frequently repeated count of colossally abundant numbers below 10^{18} , attributed to Briggs, *Experiment. Math.* **15** (2006), 251–256. That figure could not be verified against the accessible text and is not repeated here. Any use of it should re-derive it.

R.10 The formalization landscape

Karatarakis and Wiedijk formalized the Gelfond–Schneider theorem in Lean 4, including a formalized Siegel lemma for algebraic integers and the auxiliary-function machinery; that development is the basis of the port used by this corpus and cited in Appendix K. Mathlib’s Lindemann–Weierstrass development supplies analytic scaffolding only. No formalization of Baker’s theorem, of the six exponentials theorem, or of linear forms in logarithms was found in any proof assistant, so the transcendence inputs isolated behind interfaces in this report would be greenfield work.

R.11 Summary of what the literature does not supply

Three gaps are unanimous across the surveyed material. No lower-bound technology handles linear or bilinear forms whose *coefficients* are themselves logarithms, which is the wall of Section 7. No counting theorem is sensitive to the arithmetic of denominators, which is what the compact-arc reformulation needs. And no unconditional algebraic-independence result reaches products of two logarithms, which is what the criterion of Section 12 would need. The conjecture remains open, and every route surveyed inherits the same rank-two-to-rank-one collapse.

Status: [classical/open background] throughout, with the two retractions marked as such. Nothing in this appendix is used as a step in any argument of this report; it records provenance and the state of the surrounding literature.

References

- [1] L. Alaoglu and P. Erdős, *On highly composite and similar numbers*, *Transactions of the American Mathematical Society* **56** (1944), 448–469. DOI: [10.1090/S0002-9947-1944-0011087-2](https://doi.org/10.1090/S0002-9947-1944-0011087-2).
- [2] S. Ramanujan, *Highly composite numbers*, *Proceedings of the London Mathematical Society* (2) **14** (1915), 347–409; annotated manuscript in *Ramanujan Journal* **1** (1997), 119–153.
- [3] A. Baker, *Transcendental Number Theory*, Cambridge University Press, 1975.

- [4] L.-K. Hua, *Introduction to Number Theory*, Springer, 1982, Chapter 17.9.
- [5] L. A. Butler, *A Diophantine approach to the three and four exponentials conjectures*, Ramanujan Journal **42**(1) (2017), 199–221. DOI: [10.1007/s11139-015-9728-2](https://doi.org/10.1007/s11139-015-9728-2).
- [6] S. Karlin and W. J. Studden, *Tchebycheff Systems: With Applications in Analysis and Statistics*, Interscience, 1966.
- [7] L. Kuipers and H. Niederreiter, *Uniform Distribution of Sequences*, Wiley, 1974; Dover reprint, 2006.
- [8] D. Roy, *Matrices whose coefficients are linear forms in logarithms*, Journal of Number Theory **41**(1) (1992), 22–47. DOI: [10.1016/0022-314X\(92\)90081-Y](https://doi.org/10.1016/0022-314X(92)90081-Y).
- [9] Th. Schneider, *Einführung in die transzendenten Zahlen*, Springer, 1957.
- [10] M. Waldschmidt, *Diophantine Approximation on Linear Algebraic Groups: Transcendence Properties of the Exponential Function in Several Variables*, Grundlehren der mathematischen Wissenschaften 326, Springer, 2000. DOI: [10.1007/978-3-662-11569-5](https://doi.org/10.1007/978-3-662-11569-5).
- [11] M. Waldschmidt, *Variations on the Six Exponentials Theorem*, in *Algebra and Number Theory*, pp. 338–355, 2005. DOI: [10.1007/978-93-86279-23-1_23](https://doi.org/10.1007/978-93-86279-23-1_23).
- [12] M. Waldschmidt, *Six exponentials theorem — irrationality*, expository note.
- [13] M. Waldschmidt, *The Four Exponentials Problem and Schanuel’s Conjecture*, in *Mathematics Going Forward*, Lecture Notes in Mathematics, 2023, pp. 579–592. DOI: [10.1007/978-3-031-12244-6_39](https://doi.org/10.1007/978-3-031-12244-6_39).
- [14] TheoremDB contributors, *Four exponentials conjecture*, problem P4, reviewed July 31, 2026. <https://theoremdb.org/statements/P4/>.
- [15] F. Johansson et al., *mpmath: a Python library for arbitrary-precision floating-point arithmetic*, version 1.3.0. <https://mpmath.org/>.
- [16] The ProveIt project, *ProveIt repository*, <https://github.com/VladimirReshetnikov/ProveIt>, pinned in this report at commit 5d3982ceb12644fdd3499569d9db5f62f9a31dc6.
- [17] The ProveIt project, *The Alaoglu–Erdős Integer-Exponent Conjecture: Machine-Checked Partial Results, Exact Conditional Structure, and Research Directions*, source directory Analysis/ExponentialIdentities/Article/alaoglu-erdos-combined-report/, file alaoglu-erdos-combined-report.tex, August 20, 2026.
- [18] The ProveIt project, *The Alaoglu–Erdős Integer-Exponent Conjecture: Further Determinant Methods*, source directory Analysis/ExponentialIdentities/Article/alaoglu-erdos-further-report/, file alaoglu-erdos-further-report.tex, August 20, 2026.
- [19] A. Baker and G. Wüstholz, *Logarithmic forms and group varieties*, Journal für die reine und angewandte Mathematik **442** (1993), 19–62. DOI: [10.1515/crll.1993.442.19](https://doi.org/10.1515/crll.1993.442.19).
- [20] E. M. Matveev, *An explicit lower bound for a homogeneous rational linear form in the logarithms of algebraic numbers. II*, Izvestiya: Mathematics **64**(6) (2000), 1217–1269. DOI: [10.1070/IM2000v064n06ABEH000314](https://doi.org/10.1070/IM2000v064n06ABEH000314).

- [21] Y. Bugeaud and M. Laurent, *Minoration effective de la distance p -adique entre puissances de nombres algébriques*, Journal of Number Theory **61**(2) (1996), 311–342. DOI: [10.1006/jnth.1996.0152](https://doi.org/10.1006/jnth.1996.0152).
- [22] Y. Bugeaud, *Linear forms in two m -adic logarithms and applications to Diophantine problems*, Compositio Mathematica **132**(2) (2002), 137–158. DOI: [10.1023/A:1015825809661](https://doi.org/10.1023/A:1015825809661).
- [23] K. Yu, *p -adic logarithmic forms and group varieties I*, Journal für die reine und angewandte Mathematik **502** (1998), 29–92. DOI: [10.1515/crll.1998.090](https://doi.org/10.1515/crll.1998.090).
- [24] Q. Wu, *On the linear independence measure of logarithms of rational numbers*, Mathematics of Computation **72**(242) (2003), 901–911. DOI: [10.1090/S0025-5718-02-01442-4](https://doi.org/10.1090/S0025-5718-02-01442-4).
- [25] M. Morse and G. A. Hedlund, *Symbolic dynamics II. Sturmian trajectories*, American Journal of Mathematics **62**(1) (1940), 1–42. DOI: [10.2307/2371431](https://doi.org/10.2307/2371431).
- [26] M. Lothaire, *Algebraic Combinatorics on Words*, Encyclopedia of Mathematics and its Applications 90, Cambridge University Press, 2002.
- [27] A. Cobham, *Uniform tag sequences*, Mathematical Systems Theory **6** (1972), 164–192. DOI: [10.1007/BF01706087](https://doi.org/10.1007/BF01706087).
- [28] J.-P. Allouche and J. Shallit, *Automatic Sequences: Theory, Applications, Generalizations*, Cambridge University Press, 2003.
- [29] Harish-Chandra, *Differential operators on a semisimple Lie algebra*, American Journal of Mathematics **79**(1) (1957), 87–120. DOI: [10.2307/2372387](https://doi.org/10.2307/2372387).
- [30] C. Itzykson and J.-B. Zuber, *The planar approximation. II*, Journal of Mathematical Physics **21**(3) (1980), 411–421. DOI: [10.1063/1.524438](https://doi.org/10.1063/1.524438).
- [31] C. McSwiggen, *The Harish–Chandra integral: an introduction with examples*, L’Enseignement Mathématique **67**(3–4) (2021), 229–299; preprint arXiv:1806.11155. DOI: [10.4171/LEM/1017](https://doi.org/10.4171/LEM/1017).
- [32] J. Pila and A. J. Wilkie, *The rational points of a definable set*, Duke Mathematical Journal **133**(3) (2006), 591–616. DOI: [10.1215/S0012-7094-06-13336-7](https://doi.org/10.1215/S0012-7094-06-13336-7).
- [33] E. Bombieri and W. Gubler, *Heights in Diophantine Geometry*, Cambridge University Press, 2006.
- [34] G. H. Hardy, J. E. Littlewood, and G. Pólya, *Inequalities*, second edition, Cambridge University Press, 1952.
- [35] M. Laurent, M. Mignotte, and Yu. Nesterenko, *Formes linéaires en deux logarithmes et déterminants d’interpolation*, Journal of Number Theory **55** (1995), 285–321.
- [36] S. Lang, *Introduction to Transcendental Numbers*, Addison–Wesley, 1966.
- [37] S. Dasgupta and M. Kakde, *Ranks of matrices of logarithms of algebraic numbers II: The Matrix Coefficient Conjecture*, Advances in Mathematics **487** (2026), Article 110753. DOI: [10.1016/j.aim.2025.110753](https://doi.org/10.1016/j.aim.2025.110753); arXiv:2408.08178.
- [38] H. Massold, *Transcendence degrees of fields generated by exponentials of products*, arXiv:2506.01123v1, June 2025.

- [39] J.-H. Evertse, H. P. Schlickewei, and W. M. Schmidt, *Linear equations in variables which lie in a multiplicative group*, *Annals of Mathematics* **155** (2002), 807–836.
- [40] F. Carlson, *"Uber Potenzreihen mit ganzzahligen Koeffizienten"*, *Mathematische Zeitschrift* **9** (1921), 1–13.
- [41] C. Lech, *A note on recurring series*, *Arkiv för Matematik* **2** (1953), 417–421.
- [42] P. B. Borwein, *On the irrationality of certain series*, *Mathematical Proceedings of the Cambridge Philosophical Society* **112** (1992), 141–146.
- [43] M. Kawashima, *Linear independence of periods related to polylogarithms*, arXiv:2605.17756v1, May 2026.
- [44] F. Calegari, V. Dimitrov, and Y. Tang, *Arithmetic holonomy bounds and effective Diophantine approximation*, *Proceedings of the International Congress of Mathematicians 2026*, vol. 3, 356–376; arXiv:2510.04156.
- [45] J. Harrison, *Additive relations in irrational powers*, arXiv:2512.04081v3, July 2026.
- [46] A. Selberg, *Bemerkninger om et multipelt integral*, *Norsk Matematisk Tidsskrift* **26** (1944), 71–78.