

The Thue–Morse Sequence: Formula Atlas and Fabius–Rvachev Frontier Results

Consolidated 28 August 2026

Contents

I	A Unified Formula Atlas for the Thue–Morse Sequence	7
1	Four-draft merge	8
1.1	Two normalizations	8
1.2	The Pascal-row formula being complemented	8
1.3	Merge policy	8
1.4	A compact formula index	9
2	Binary character and automata	10
2.1	Recursion and uniqueness	10
2.2	Binary concatenation and reflection	11
2.3	Exclusive-or is the natural group law	12
2.4	Finite automata, matrices, and tensor powers	12
2.5	The sequence is primitive recursive	13
2.6	Exact weight distribution in a dyadic block	14
2.7	Evil and odious enumerators	15
3	Digit and trigonometric formulae	15
3.1	Extracting one binary digit	15
3.2	Dyadic floor and fractional-part sums	16
3.3	Square waves and a sine-sign product	17
3.4	The sign of the Fabius sinc product	18
3.5	The ruler transition	21
4	Valuations and carries	21
4.1	Legendre’s formula	21
4.2	A single central binomial coefficient	22
4.3	Catalan valuations	23
4.4	Kummer’s carry cocycle	23
4.5	Subtraction and borrow counts	24
5	Pascal support and Boolean inversion	24
5.1	Lucas’s theorem as a submask criterion	24
5.2	A logarithm-free Pascal-row formula modulo three	25
5.3	A one-parameter row statistic	26
5.4	Sparse Pascal probes for individual bits	26

5.5	Odd multinomial coefficients and exact r -nomial logarithms	27
5.6	Boolean weight enumerators and Möbius inversion	28
6	Generating products	30
6.1	The Boolean-cube mother identity	30
6.2	The finite Prouhet polynomial	30
6.3	Reciprocity and cyclotomic multiplicities	32
6.4	The infinite product and Mahler equation	33
6.5	A weighted digit enumerator	35
6.6	Thue–Morse constants in every base	36
7	Euler transforms and algebraic series	36
7.1	The logarithmic ruler transform	36
7.2	A nonlocal convolution recurrence	37
7.3	Partition and Bell-polynomial formulae	38
7.4	A Hessenberg determinant	39
7.5	The zero–one series over $\mathbb{F}_2$	40
7.6	An integral algebraic lift	41
8	Natural boundary and Hankel rigidity	42
8.1	The unit circle is a natural boundary	42
8.2	Hankel determinants never vanish	43
8.3	Transcendental values and the sharp irrationality exponent	44
9	Finite differences and moments	45
9.1	The mixed-difference master identity	45
9.2	Exact multiple-integral remainder	47
9.3	Arbitrary translation weights	48
9.4	Prouhet cancellation and the first surviving moment	49
9.5	The finite exponential transform	50
9.6	All power moments: a finite composition formula	52
9.7	Bell-polynomial compression and cumulants	52
9.8	Centered moments and a parity selection rule	53
9.9	All binomial-basis moments	54
10	Sparse Prouhet identities	55
11	Prefix sums and the Fabius bridge	57
11.1	The first prefix discrepancy	57
11.2	Iterated inclusive prefixes	58
11.3	Terminal zero runs at dyadic boundaries	59
11.4	Why this is the discrete spine of the Fabius theory	60
12	Fourier transforms	60
12.1	A finite Fourier product	60
12.2	Sine and cosine products as Thue–Morse sums	61
12.3	The exact dyadic discrete Fourier transform	62

13 Walsh, Riesz, and autocorrelation	63
13.1 A one-point Walsh spectrum	63
13.2 Rademacher step functions	63
13.3 Finite Riesz products	64
13.4 Finite and limiting autocorrelation	65
13.5 A Mahler equation for the autocorrelation series	66
14 Word combinatorics and dynamics	67
14.1 Uniform recurrence and letter frequencies	67
14.2 Aperiodicity	68
14.3 Overlap-freeness and the critical exponent	68
14.4 Complement and reversal closure	69
14.5 Linear factor complexity	69
15 Rarefied sums and Newman’s phenomenon	71
15.1 Root-of-unity filters at dyadic endpoints	71
15.2 Complete evaluation modulo three	71
15.3 Gelfond’s uniform exponent	72
15.4 A matrix recursion for odd moduli	73
15.5 The global Newman theorem and the arithmetic frontier	74
16 Mellin–Dirichlet continuation	75
16.1 Quadratic logarithmic decay near one	75
16.2 The shifted Dirichlet family	76
16.3 Mellin representation and entire continuation	77
16.4 The harmonic kernel	78
17 Thue–Morse exponent products	78
17.1 Pair reduction	78
17.2 A two-parameter master product	79
17.3 The Woods–Robbins product	80
17.4 Three dyadic block-product limits	81
18 Base-q generalization	82
19 Formula selection and Lean roadmap	84
19.1 Formalized modules (August 2026)	84
19.2 Which representation is best?	93
19.3 A compact verification script	94
19.4 A dependency-aware formalization order	94
19.5 Proof-dependency graph	97
20 Worked examples	97
20.1 One index through many representations	97
20.2 An all-order dyadic moment	98
20.3 A Fourier sample	98
20.4 Autocorrelation from the recursion	98
20.5 A complete four-bit block	99
20.6 A carry-corrected addition	99
21 Conclusion	99

A	First 64 terms	100
B	Identity dictionary	100
C	Draft concordance	101
II	Finite-Block Calculus for the Fabius–Rvachev–Thue–Morse System	104
1	Corpus audit and research positioning	105
1.1	What was inspected	105
1.2	Established pillars	105
1.3	The missing commutative bridge	106
1.4	Priority and uncertainty	106
2	Objects, normalizations, and exact background	106
2.1	The signed block polynomial	106
2.2	Fourier and Laplace products	107
2.3	The lower-Lambert phase	107
3	One Thue–Morse block, two exact product bridges	108
3.1	Oscillatory specialization: the Fourier bridge	108
3.2	Decaying specialization: the endpoint Laplace bridge	108
3.3	The omitted tail is a scaled copy of the whole	109
4	The zeta–Lambert tail calculus	110
4.1	A local logarithm of the sinc function	110
4.2	All orders, exact radius, explicit remainder	110
4.3	Arbitrary-order finite Thue–Morse approximants	112
4.4	Why the coefficient array is a Lambert series	113
5	An all-orders finite-block formula for the endpoint Laplace product	113
5.1	Centering removes the linear tail	113
5.2	Bernoulli–zeta expansion of the negative tail	114
5.3	The endpoint product from one finite digital block	114
6	q-binomial Richardson acceleration on the spectral scale	115
6.1	An analytic error germ in q^m	115
6.2	Geometric-grid weights are Gaussian binomials	115
7	Exact random tails and a weak differential expansion	118
7.1	The random-series model	118
7.2	Cumulants are the same zeta–Lambert coefficients	119
7.3	A certified weak expansion	119
8	Lambert-W complexity of endpoint product truncation	120
8.1	The canonical scale and the tail variable	120
8.2	A finite-block endpoint evaluation protocol	121
9	A Thue–Morse Barnes-type tower from Dirichlet trivial zeros	122
9.1	Input from the Atlas	122
9.2	The derivative tower	122

10 Numerical verification and order diagnostics	124
10.1 Exact-identity residuals	124
10.2 Observed convergence orders	124
10.3 Weak expansion check	125
10.4 Endpoint-tail coefficients	125
10.5 Lambert-depth experiment	126
11 Consequences for the existing research program	127
11.1 A new exact-arithmetic/Fourier bridge	127
11.2 A new q -binomial role	127
11.3 A certified route from finite splines to up	127
11.4 A product engine for endpoint saddles	127
11.5 Formalization leverage	128
12 Lean formalization roadmap	128
13 Frontier conjectures and next deductions	132
13.1 General radix and nonuniform tails	133
14 Conclusion	133
A Coefficient and extrapolation tables	135
A.1 Zeta–Lambert and centered Laplace coefficients	135
A.2 Analytic error germ	135
A.3 Gaussian-binomial extrapolation weights	135
B Reproducibility protocol	136
C Complete corpus ledger and deduplication map	137
C.1 Documentation treated as authoritative at the audit snapshot	137
C.2 Frozen archive represented through provenance	138

Provenance

This volume preserves the absorbed theorem and proof bodies, with labels, citation keys, and asset paths mechanically prefixed per part and wrapper metadata and counter state normalized. Later edits add status-only Lean crosswalk annotations; the historical source hashes below are unchanged. The absorbed drafts are deleted from the working tree; git history is the archive.

- Thue_Morse_Formula_Atlas/Thue_Morse_Formula_Atlas.tex
SHA-256 bafeee49431ae8d65e16ec0b978ca4adda6b10b82dbe2c0ddf10def444334795
- Fabius_Rvachev_Thue_Morse_Frontier_Results/Fabius_Rvachev_Thue_Morse_Frontier_Results.tex
SHA-256 3afa1618db2b7f67e0ec2652ad6cea3fb9429b60fa0391494eb32035d34e3063

Part I

A Unified Formula Atlas for the Thue–Morse Sequence

Abstract

Four draft reports in the Fabius-function research-frontier directory developed partly overlapping collections of formulae for the Thue–Morse sequence. This document merges them into one proof-oriented atlas, removes duplicated arguments, reconciles notation and sign conventions, and retains the strongest version of each genuinely different result. The signed sequence $\varepsilon_n = \varepsilon_n$ and the zero–one sequence $\tau(n) = \text{wt}_2(n) \bmod 2$ are treated simultaneously.

The resulting network begins with binary recursion, block concatenation, XOR characters, floor and fractional-part sums, trigonometric signs, factorial and central-binomial valuations, Catalan valuations, and Kummer carry cocycles. It then develops the Pascal-row viewpoint in several directions: Glaisher’s odd entry count, the existing exact binomial–logarithm formula, a logarithm-free modulo-three formula, sparse power-of-two probes, Boolean Möbius inversion, multinomial support counts, and a sparse Prouhet theorem on the odd support of an arbitrary Pascal row.

The lacunary product

$$\sum_{n \geq 0} \varepsilon_n z^n = \prod_{j \geq 0} (1 - z^{2^j})$$

serves as the central analytic engine. From it we obtain Mahler equations, cyclotomic factorizations, a ruler-function Euler transform, convolution, partition, Bell-polynomial, and Hessenberg-determinant formulae, an algebraic formal series over $\mathbb{F}_2$, an integral algebraic lift, exact finite differences, all dyadic power moments, Walsh and ordinary Fourier transforms, Riesz products, autocorrelation recurrences, iterated prefix sums, and the bridge back to the Fabius and Rvachev functions.

New material added in the unification includes an explicit dyadic weight-distribution chapter, a systematic centered-moment parity theorem, and a Mellin–Dirichlet theory. In particular, the shifted Dirichlet series

$$D(s, a) = \sum_{n \geq 0} \frac{\varepsilon_n}{(n + a)^s}$$

continues to an entire function, satisfies an exact dyadic parameter equation, and vanishes at every nonpositive integer. The proof is driven by the super-polynomial flatness

$$\log \prod_{j \geq 0} (1 - e^{-2^j t}) = -\frac{\log^2(1/t)}{2 \log 2} + \mathcal{O}(\log(1/t)) \quad (t \downarrow 0).$$

A second consolidation pass merged three parallel sibling atlases into this document. It contributed the weighted digit enumerator $\prod_j (1 + uz^{2^j})$ and the subtraction/borrow cocycle; self-contained proofs that the unit circle is a natural boundary and that the Hankel determinants never vanish, with the sharp irrationality exponent $\mu(\Theta_b) = 2$ of the Thue–Morse–Mahler numbers stated from the literature; a Mahler functional equation for the autocorrelation series; a chapter on the combinatorics and dynamics of the infinite word (uniform recurrence, aperiodicity, Thue’s overlap theorem, critical exponent two, linear factor complexity); and a chapter on rarefied sums – exact residue-class sums modulo three, Gelfond’s optimal exponent $\log_4 3$, a matrix recursion for odd moduli, and Newman’s positivity phenomenon with its modern extensions along primes, squares, and cubes.

Every elementary assertion is proved in the text; deeper historical results about transcendence, substitution spectra, and evaluated products are clearly identified and cited. No mathematical priority is claimed for identities merely derived here, and no formula is silently represented as already formalized in Lean.

1 Scope, conventions, and the four-draft merge

1.1 Two normalizations

Write the binary expansion of $n \in \mathbb{N}_0 = \{0, 1, 2, \dots\}$ as

$$n = \sum_{j \geq 0} b_j(n) 2^j, \quad b_j(n) \in \{0, 1\}, \quad (1.1)$$

with only finitely many nonzero digits. Its binary weight and the two standard normalizations of the Thue–Morse sequence are

$$\text{wt}_2(n) = \sum_{j \geq 0} b_j(n), \quad \tau(n) = \text{wt}_2(n) \bmod 2 \in \{0, 1\}, \quad \varepsilon_n = \varepsilon_n = 1 - 2\tau(n). \quad (1.2)$$

Thus

$$\tau(n) = \frac{1 - \varepsilon_n}{2}, \quad \varepsilon_n = 1 - 2\tau(n). \quad (1.3)$$

The signed normalization is usually cleaner for products, transforms, and cancellation; the zero–one normalization is natural for automata, base expansions, and the formula in the source exposition.

We use $a \oplus b$ for bitwise exclusive-or and $k \preceq_2 n$ when every one-bit of k is also a one-bit of n . Empty products are 1, empty sums are 0, and $0^0 = 1$ when it occurs in a finite moment identity. The valuation $\nu_2(m)$ is the exponent of 2 in a positive integer m , with $\nu_2(1) = 0$.

1.2 The Pascal-row formula being complemented

Section 11.8 of the main Fabius–Rvachev exposition defines

$$\mathcal{X}(n) = \sum_{k=0}^n (-1)^{\binom{n}{k}} \quad (1.4)$$

and proves the exact-power identity

$$n + 1 - \mathcal{X}(n) = 2^{\text{wt}_2(n)+1}. \quad (1.5)$$

Consequently

$$\tau(n) = \frac{1 + (-1)^{\log_2 \left(n+1 - \sum_{k=0}^n (-1)^{\binom{n}{k}} \right)}}{2}. \quad (1.6)$$

The logarithm is exact because its argument was first proved to be a power of two. The formula hides the binary expansion inside the parity pattern of a whole Pascal row. The present atlas asks what other structures carry the same parity character.

1.3 Merge policy

The four input drafts were not independent papers: large parts shared the same binary, valuation, generating-product, Prouhet, and Fourier skeleton. A literal concatenation would repeat the same proof several times while obscuring the interesting differences. The unified text therefore follows three rules.

1. A theorem common to several drafts appears once, in the strongest useful form, with one proof and cross-references to its later specializations.

2. Material unique to a draft is retained when it is mathematically distinct: in particular the modulo-three Pascal formula, multinomial logarithms, sparse Pascal Prouhet identities, Boolean Möbius inversion, determinant formula, integer algebraic lift, odious/evil enumerators, autocorrelation, Dirichlet integrals, and the explicit Fabius-prefix bridge.
3. Statements added during the merge are marked by context rather than by novelty rhetoric. The report claims derivation, not historical priority.

This unified atlas was itself produced in parallel: three sibling consolidations of the same four drafts were written independently, and a second merge pass (26 August 2026) folded their distinctive material into the present document under the same three rules. The additions are listed in the abstract and in the concordance of [section C](#).

Proof-status convention. The repository contains a substantial formal development of this atlas, but a paper proof is not by itself evidence of formalization. Statements proved in Lean carry their exact declaration names in the text, typeset as `Fabius.thueMorseSign`; statements without such a citation must not be inferred to have a formal counterpart.

The current inventory in [section 19.1](#) consists of one hundred and three compiled Lean modules. It covers the finite algebraic, automata, moment, Fourier, rarefaction, Dirichlet–Mellin, word-combinatorial, and Fabius–Rvachev bridges recorded there, often in greater generality than the displayed prose. The dependency-aware map in [section 19](#) distinguishes this delivered inventory from the historical module proposal retained for comparison.

1.4 A compact formula index

The table is a navigation device rather than a substitute for hypotheses and proofs. Here $P_m(z) = \sum_{n < 2^m} \varepsilon_n z^n$, and $E(z) = \sum_{n \geq 0} \varepsilon_n z^n$.

Viewpoint	Representative exact formula
Binary recursion	$\varepsilon_{2n} = \varepsilon_n, \varepsilon_{2n+1} = -\varepsilon_n$.
Block concatenation	$\varepsilon_{2^m q + r} = \varepsilon_q \varepsilon_r$ for $0 \leq r < 2^m$.
XOR character	$\varepsilon_{a \oplus b} = \varepsilon_a \varepsilon_b$.
Digit product	$\varepsilon_n = \prod_j (1 - 2b_j(n))$.
Floors	$\varepsilon_n = (-1)^{\sum_{j \geq 0} \lfloor n/2^j \rfloor}$.
Fractional parts	$\text{wt}_2(n) = \sum_{j \geq 1} \{n/2^j\}$.
Factorials	$\varepsilon_n = (-1)^{n - \nu_2(n!)}$.
Central binomial	$\varepsilon_n = (-1)^{\nu_2 \binom{2n}{n}}$.
Carries	$\varepsilon_{a+b} = \varepsilon_a \varepsilon_b (-1)^{\nu_2 \binom{a+b}{a}}$.
Pascal support	$\#\{k : \binom{n}{k} \text{ odd}\} = 2^{\text{wt}_2(n)}$.
Sparse Pascal probes	$\varepsilon_n = (-1)^{\sum_j \binom{n}{2^j}}$.
Log-free Pascal residue	$\tau(n) = (2^{\text{wt}_2(n)} \bmod 3) - 1$.
Multinomial support	$O_r(n) = r^{\text{wt}_2(n)}$.
Boolean Möbius	$\mu_2(k, n) = \varepsilon_{n \oplus k}$ for $k \preceq_2 n$.
Finite product	$P_m(z) = \prod_{j=0}^{m-1} (1 - z^{2^j})$.
Weighted complex block	$\sum_{n < 2^m} u^{\text{wt}_2(n)} e^{a+nz} = e^a \prod_{j < m} (1 + u e^{2^j z})$, including the base case $a = 0$.

Viewpoint	Representative exact formula
Signed complex block	$\sum_{n < 2^m} \varepsilon_n e^{a+nz} = e^a \prod_{j < m} (1 - e^{2^j z})$, including the base case $a = 0$.
Infinite product	$E(z) = \prod_{j \geq 0} (1 - z^{2^j})$, $ z < 1$.
Mahler equation	$E(z) = (1 - z)E(z^2)$.
Euler transform	$\log E(z) = -\sum_{r \geq 1} (2^{\nu_2(r)+1} - 1) z^r / r$.
Finite differences	$\sum_{n < 2^m} \varepsilon_n f(x + n) = (-1)^m \Delta_1 \cdots \Delta_{2^m-1} f(x)$.
Prouhet	$\sum_{n < 2^m} \varepsilon_n p(n) = 0$ when $\deg p < m$.
Sharp moment	$\sum_{n < 2^m} \varepsilon_n n^m = (-1)^m m! 2^{\binom{m}{2}}$.
Sparse Prouhet	$\sum_{k \preceq_2 n} \varepsilon_k p(k) = 0$ when $\deg p < \text{wt}_2(n)$.
Prefix discrepancy	$\sum_{n < N} \varepsilon_n = 0$ for even N , $\varepsilon_{(N-1)/2}$ for odd N .
Walsh transform	A length- 2^m block is one Walsh frequency.
Ordinary DFT	$\text{DFT}_{2^m}[\varepsilon][k] = \prod_{j < m} (1 - e^{-2\pi i k 2^j / 2^m})$.
Riesz density	$2^{-m} P_m(e^{2\pi i x}) ^2 = \prod_{j < m} (1 - \cos(2\pi 2^j x))$.
Characteristic two	$(1+x)^3 T(x)^2 + (1+x)^2 T(x) + x = 0$ in $\mathbb{F}_2[[x]]$.
Weighted enumerator	$\sum_n u^{\text{wt}_2(n)} z^n = \prod_j (1 + u z^{2^j})$.
Subtraction cocycle	$\varepsilon_{a-b} = \varepsilon_a \varepsilon_b (-1)^{\nu_2 \binom{a}{b}}$ for $b \leq a$.
Autocorrelation Mahler	$2zC(z) = 1 - (1-z)^2 C(z^2)$.
Natural boundary	E continues across no arc of $ z = 1$.
Hankel rigidity	$\det(\varepsilon_{i+j})_{i,j < r} \neq 0$ for every r .
Rarefaction mod 3	$(A_{2u,r}^{(3)})_r = 3^{u-1} (2, -1, -1)$.
Gelfond exponent	$3^{m/2} \leq \sup_x P_m(e^{2\pi i x}) \leq \frac{2}{\sqrt{3}} 3^{m/2}$.
Newman positivity	$\sum_{n < X, 3 n} \varepsilon_n > 0$ for all $X \geq 1$.
Critical exponent	Overlap-free; squares occur, exponent exactly 2.
Factor complexity	$\ell + 1 \leq p(\ell) \leq 6\ell - 4$ for $\ell \geq 1$.
Dirichlet–Mellin	$D(s, a) = \Gamma(s)^{-1} \int_0^\infty t^{s-1} e^{-at} E(e^{-t}) dt$. <i>Delivered</i> : boundary flatness, convergence, the dyadic equation, the analytic infinite product, the Mellin representation, and the entire continuation (ThueMorseBoundaryFlatness.lean, ThueMorseDirichlet.lean, ThueMorseInfiniteProduct.lean, ThueMorseMellin.lean, ThueMorseEntireContinuation.lean).

2 Binary character, blocks, automata, and enumerators

2.1 Recursion and uniqueness

Theorem 2.1 (Binary recursion and uniqueness). *The signed sequence satisfies*

$$\varepsilon_0 = 1, \quad \varepsilon_{2n} = \varepsilon_n, \quad \varepsilon_{2n+1} = -\varepsilon_n. \quad (2.1)$$

Equivalently,

$$\tau(0) = 0, \quad \tau(2n) = \tau(n), \quad \tau(2n+1) = 1 - \tau(n). \quad (2.2)$$

Either system and its initial value determines the sequence uniquely.

Proof. Appending a terminal binary digit 0 leaves the digit sum unchanged; appending a terminal 1 increases it by one. For uniqueness, every positive integer is $2q$ or $2q + 1$ with $q < n$, so strong induction eventually reduces its value to the initial value at zero. $\square$

Let $W_m = \tau(0)\tau(1) \cdots \tau(2^m - 1)$, and let $\overline{W}$ denote bitwise complement. The recurrence gives

$$W_{m+1} = W_m \overline{W_m}. \quad (2.3)$$

Thus the infinite word is the fixed point beginning with 0 of

$$0 \mapsto 01, \quad 1 \mapsto 10. \quad (2.4)$$

In signed form a block is followed by its negative.

Formal versions. The signed recursion (2.1) is `Fabius.thueMorseSign_two_mul` and `Fabius.thueMorseSign_two_mul_add_one` in `DyadicClosedForm.lean`, while the bit form (2.2) is `Fabius.thueMorseBit_two_mul` and `Fabius.thueMorseBit_two_mul_add_one` in `ThueMorseEnumerators.lean`. The block-complement law (2.3), in the signed reading stated here — a block is followed by its negative — is `Fabius.thueMorseSign_add_pow_two`, $\varepsilon_{h+2^b} = -\varepsilon_h$ for $h < 2^b$, in the same file. The uniqueness clause is not separately formalized: in Lean the sequence is *defined* by its binary weight and the recursions are theorems about it, so there is no second candidate to rule out.

2.2 Binary concatenation and reflection

Theorem 2.2 (Block concatenation). *For $m, q \geq 0$ and $0 \leq r < 2^m$,*

$$\text{wt}_2(2^m q + r) = \text{wt}_2(q) + \text{wt}_2(r), \quad (2.5)$$

$$\varepsilon_{2^m q + r} = \varepsilon_q \varepsilon_r, \quad (2.6)$$

$$\tau(2^m q + r) = \tau(q) \oplus \tau(r). \quad (2.7)$$

Proof. Multiplication by 2^m shifts the binary expansion of q by m places. Since $r < 2^m$, adding r fills only those vacated positions and creates no carry. The bits concatenate, so weights add and parities follow. $\square$

Corollary 2.3 (Reflection and bit-position symmetry). *For $0 \leq n < 2^m$,*

$$\varepsilon_{2^m - 1 - n} = (-1)^m \varepsilon_n, \quad \tau(2^m - 1 - n) = \tau(n) \oplus (m \bmod 2). \quad (2.8)$$

Permuting the first m binary positions leaves ε_n and $\tau(n)$ unchanged.

Proof. Within m positions, $2^m - 1 - n$ is the bitwise complement of n , so its weight is $m - \text{wt}_2(n)$. A permutation of positions leaves their sum unchanged. $\square$

Formal versions. The signed reflection (2.8) is `Fabius.thueMorseSign_dyadic_complement` (`DyadicClosedForm.lean`), proved by induction on m through the two-scale laws rather than through the weight of the complement; its only hypothesis is the stated $n < 2^m$. Only the signed form is named in Lean; the bit form is its parity reading, and the digit-permutation invariance is the invariance of wt_2 under permuting positions.

2.3 Exclusive-or is the natural group law

Theorem 2.4 (XOR character). *For all $a, b \geq 0$,*

$$\tau(a \oplus b) = \tau(a) \oplus \tau(b), \quad \boxed{\varepsilon_{a \oplus b} = \varepsilon_a \varepsilon_b.} \quad (2.9)$$

Thus $n \mapsto \varepsilon_n$ is a character of the Boolean group $\bigoplus_{j \geq 0} \mathbb{Z}/2\mathbb{Z}$.

Proof. At each position, XOR adds the two bits modulo two. Summing the coordinate identities modulo two gives $\text{wt}_2(a \oplus b) \equiv \text{wt}_2(a) + \text{wt}_2(b) \pmod{2}$. $\square$

Corollary 2.5 (XOR autocorrelation and Hamming spheres). *Fix $m \geq 0$, $0 \leq a < 2^m$, and $0 \leq r \leq m$. Then*

$$\sum_{n=0}^{2^m-1} \varepsilon_n \varepsilon_{n \oplus a} = 2^m \varepsilon_a, \quad (2.10)$$

$$\sum_{\substack{0 \leq n < 2^m \\ \text{wt}_2(n \oplus a) = r}} \varepsilon_n = \varepsilon_a (-1)^r \binom{m}{r}. \quad (2.11)$$

Proof. The first summand is termwise ε_a . For the second, put $h = n \oplus a$. Then $\varepsilon_n = \varepsilon_a \varepsilon_h = \varepsilon_a (-1)^r$, and there are $\binom{m}{r}$ masks of weight r . $\square$

Formal versions. `Fabius.sum_thueMorseSign_mul_xor` (the autocorrelation, which the Lean proof establishes termwise from the character identity, with no alignment hypothesis on a) and `Fabius.sum_thueMorseSign_hammingSphere` (the sphere sums, by reindexing along the XOR involution), in `ThueMorseEnumerators.lean`.

2.4 Finite automata, matrices, and tensor powers

The recurrence gives a two-state automaton: begin in state $+1$, read the binary digits in either order, and flip the state for each one-bit. In matrix form, let

$$M_0 = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}, \quad M_1 = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}, \quad e_0 = \begin{pmatrix} 1 \\ 0 \end{pmatrix}. \quad (2.12)$$

If $n = \sum_{j=0}^{m-1} b_j 2^j$, then

$$M_{b_0} \cdots M_{b_{m-1}} e_0 = \begin{pmatrix} 1 - \tau(n) \\ \tau(n) \end{pmatrix}. \quad (2.13)$$

The matrices commute, reflecting the fact that only the number of one-bits matters. In natural binary order,

$$(\varepsilon_0, \dots, \varepsilon_{2^m-1})^\top = ((1, -1)^\top)^{\otimes m}. \quad (2.14)$$

Hence every residue subsequence in the 2-kernel is either ε or $-\varepsilon$:

$$(\varepsilon_{2^m n + r})_{n \geq 0} = \varepsilon_r (\varepsilon_n)_{n \geq 0}. \quad (2.15)$$

This is the finite-kernel criterion for 2-automaticity in its smallest possible form.

Algorithm 2.6 (XOR-fold evaluator). For a fixed-width machine word, fold the bits by XORing with successively larger right shifts. The final least significant bit is $\tau(n)$.

```

x := n
x := x xor (x >> 1)
x := x xor (x >> 2)
x := x xor (x >> 4)
x := x xor (x >> 8)
... continue while the shift is below the word width ...
tau := x & 1
epsilon := 1 - 2*tau

```

2.5 The sequence is primitive recursive

The automaton above bounds the *memory* needed to generate the sequence. The following statement bounds the *recursion* needed to define it, and it is what licenses the fixed, computable step count of Algorithm 2.6.

Theorem 2.7 (Primitive recursiveness). *The binary weight wt_2 , the zero–one sequence τ , and the signed sequence ε are primitive recursive functions. Consequently all three are computable, and the evil and odious sets are decidable.*

This is not quite the content of (2.2). That recurrence computes $\tau(n)$ from $\tau(\lfloor n/2 \rfloor)$, whereas the official definition of a primitive recursive function permits a recursive call only at the immediate predecessor $n - 1$. A recursion that descends by halving is a *course-of-values* recursion, and the device that reduces it to the official scheme is to carry the entire history along. We isolate that device once, in the generality every base- b digit recursion needs.

Lemma 2.8 (Recursion along a primitive recursive descent). *Let d be primitive recursive with $d(n) < n$ for every $n \geq 1$, let g be primitive recursive, and let z be a constant. Then the unique function f satisfying*

$$f(0) = z, \quad f(n) = g(n, f(d(n))) \quad (n \geq 1) \quad (2.16)$$

is primitive recursive. No hypothesis is imposed on $d(0)$.

Proof. Let $F(n)$ code the list $\langle f(0), \dots, f(n-1) \rangle$ of all values below n , and let $\lambda(L, i)$ be the primitive recursive lookup that returns the i -th entry of the list coded by L when i is in range, and reports failure otherwise. Then F obeys an *ordinary* primitive recursion,

$$F(0) = \langle \rangle, \quad F(n+1) = F(n) \frown \begin{cases} g(n, \lambda(F(n), d(n))) & \text{if that lookup succeeds,} \\ z & \text{otherwise.} \end{cases}$$

For $n \geq 1$ the entry demanded carries index $d(n) < n$ and is therefore already present in the length- n list $F(n)$, so the appended value is $g(n, f(d(n))) = f(n)$; at $n = 0$ the list is empty, every lookup fails, and the appended value is $z = f(0)$. Note that the fallback must guard the *whole* step expression rather than the lookup alone — defaulting only the lookup would append $g(0, z)$ instead of z — and that this is exactly what frees the statement from any assumption on $d(0)$. Hence F is primitive recursive, and so is $f(n) = \lambda(F(n+1), n)$. $\square$

Proof of Theorem 2.7. Deleting the last binary digit of n deletes exactly that digit from the digit sum, so

$$\text{wt}_2(n) = \text{wt}_2(\lfloor n/2 \rfloor) + (n \bmod 2), \quad (2.17)$$

an identity that holds at $n = 0$ as well, both sides being zero. This is (2.16) with $d(n) = \lfloor n/2 \rfloor$, which is primitive recursive and satisfies $d(n) < n$ for $n \geq 1$; with $z = 0$; and with $g(n, s) = s + (n \bmod 2)$, primitive recursive because addition and remainder are. Lemma 2.8 then makes wt_2 primitive recursive, whence $\tau = \text{wt}_2 \bmod 2$ is primitive recursive by composition, and so is $\varepsilon = 1 - 2\tau$ by (1.2). Primitive recursive functions are total and computable, and a set whose characteristic function is primitive recursive is decidable. $\square$

Remark 2.9. Identity (2.17) is the unsigned form of (2.2): reducing it modulo two gives $\tau(n) = \tau(\lfloor n/2 \rfloor) \oplus (n \bmod 2)$, which is the two clauses $\tau(2n) = \tau(n)$ and $\tau(2n+1) = 1 - \tau(n)$ written as one. The bounded expression $\text{wt}_2(n) = \sum_{i < n} (\lfloor n/2^i \rfloor \bmod 2)$ yields a second, recursion-free proof, because bounded sums of primitive recursive functions are primitive recursive; the descent formulation is preferred here because it is the one that survives the passage to base q in Section 18.

Formal versions. Theorem 2.7 and Lemma 2.8 are fully formal (`ThueMorseComputability.lean`). The descent scheme is `Fabius.primrec_of_descent`, stated for an arbitrary primitive recursive d with $d(n) < n$ and an arbitrary `Primcodable` codomain — more general than the halving case needed here, which is recorded separately as `Fabius.primrec_of_halving`. The halving identity (2.17) is `Fabius.binaryWeight_div_two`, and the three primitive-recursiveness statements are `Fabius.primrec_binaryWeight`, `Fabius.primrec_thueMorseBit`, and `Fabius.primrec_thueMorseSign`; `Fabius.nat_primrec_binaryWeight` and `Fabius.nat_primrec_thueMorseBit` record the same facts in the textbook form `Nat.Primrec` for functions $\mathbb{N}_0 \rightarrow \mathbb{N}_0$. The computability corollaries are `Fabius.computable_thueMorseBit` and `Fabius.computable_thueMorseSign`, and decidability of the odious set is `Fabius.primrecPred_thueMorseBool`. The signed sequence needs no integer arithmetic at all: every function out of `Bool` is primitive recursive, so ε is the Boolean word followed by $b \mapsto \pm 1$. The Boolean word itself (`Fabius.thueMorseBool`) carries the automaticity clauses `Fabius.thueMorseBool_two_mul` and `Fabius.thueMorseBool_two_mul_add_one`. Finally `Fabius.primrec_thueMorsePrefix` states that $n \mapsto \tau(0)\tau(1)\cdots\tau(n-1)$ is primitive recursive: the infinite word can be *generated*, not merely decided one letter at a time.

2.6 Exact weight distribution in a dyadic block

The four drafts repeatedly use the fact that a length- 2^m block is a Boolean cube. It is useful to record the complete distribution, not only its parity consequence.

Theorem 2.10 (Binomial weight enumerator). *For every $m \geq 0$,*

$$\sum_{n=0}^{2^m-1} u^{\text{wt}_2(n)} = (1+u)^m. \quad (2.18)$$

Consequently

$$\#\{0 \leq n < 2^m : \text{wt}_2(n) = r\} = \binom{m}{r}. \quad (2.19)$$

Under the uniform distribution on the block, $\text{wt}_2(n)$ is binomial with mean $m/2$ and variance $m/4$.

Proof. Each of the m positions independently contributes 1 or u . Multiplication gives $(1+u)^m$; coefficient extraction and the standard binomial moments give the remaining statements. $\square$

Formal versions. The enumerator is `Fabius.sum_pow_binaryWeight_eq_one_add_pow`, proved over an arbitrary commutative semiring, and the count (2.19) is `Fabius.card_filter_binaryWeight_eq`, extracted from it exactly as in the proof above: read the enumerator in $\mathbb{N}_0[z]$ and compare the coefficient of z^r . Both live in `ThueMorseEnumerators.lean`; the underlying Boolean-cube machinery is described after theorem 6.2.

Setting $u = -1$ gives exact sign balance for $m \geq 1$. More generally, the characteristic function of the weight is

$$2^{-m} \sum_{n < 2^m} e^{it \text{wt}_2(n)} = \left(\frac{1 + e^{it}}{2} \right)^m. \quad (2.20)$$

This elementary statistical layer is useful when comparing Boolean/Walsh structure with ordinary Fourier structure later in the report.

2.7 Evil and odious enumerators

Call n *evil* when $\text{wt}_2(n)$ is even and *odious* when it is odd. Let $\mathcal{E}_n$ and $\mathcal{O}_n$ be the increasing enumerations, starting at index zero.

Theorem 2.11 (Evil and odious numbers). *For every $n \geq 0$,*

$$\boxed{\mathcal{E}_n = 2n + \tau(n), \quad \mathcal{O}_n = 2n + 1 - \tau(n).} \quad (2.21)$$

Hence

$$\mathcal{O}_n - \mathcal{E}_n = \varepsilon_n, \quad \mathcal{O}_n + \mathcal{E}_n = 4n + 1. \quad (2.22)$$

Proof. The two consecutive integers $2n$ and $2n + 1$ have parities $\tau(n)$ and $1 - \tau(n)$ of their binary weights. If $\tau(n) = 0$, then $2n$ is evil and $2n + 1$ odious; if $\tau(n) = 1$, the roles reverse. The displayed formulas choose the correct member of each consecutive pair. $\square$

Corollary 2.12 (Composition laws). *For every $n \geq 0$,*

$$\mathcal{E}_{\mathcal{E}_n} = 2\mathcal{E}_n, \quad \mathcal{O}_{\mathcal{E}_n} = 2\mathcal{E}_n + 1, \quad (2.23)$$

$$\mathcal{E}_{\mathcal{O}_n} = 2\mathcal{O}_n + 1, \quad \mathcal{O}_{\mathcal{O}_n} = 2\mathcal{O}_n. \quad (2.24)$$

Proof. By definition $\tau(\mathcal{E}_n) = 0$ and $\tau(\mathcal{O}_n) = 1$. Substitute these values into (2.21) with the new index. $\square$

Formal versions. In `ThueMorseEnumerators.lean` the enumerations are `Fabius.evileNum` and `Fabius.odiousEnum`, with landing lemmas `Fabius.thueMorseBit_evileNum`, `Fabius.thueMorseBit_odiousEnum`, strict monotonicity `Fabius.evileNum_strictMono`, `Fabius.odiousEnum_strictMono`, difference and sum `Fabius.odiousEnum_sub_evileNum`, `Fabius.evileNum_add_odiousEnum`, and the four composition laws. The Lean proof of surjectivity isolates a detail worth recording in prose: for every m with $\tau(m) = 0$ one has $\mathcal{E}_{\lfloor m/2 \rfloor} = m$ (`Fabius.evileNum_div_two`) – the same index $\lfloor m/2 \rfloor$ works whether m is even (then $\tau(m/2) = 0$ and the τ -summand vanishes) or odd (then $\tau((m-1)/2) = 1$ and the summand supplies the missing one) – and dually for odious m (`Fabius.odiousEnum_div_two`). Monotone enumerations that hit everything are complete, so no separate counting argument is needed.

3 Digit, floor, fractional-part, and trigonometric formulas

3.1 Extracting one binary digit

Lemma 3.1 (Floor extraction). *For every $j, n \geq 0$,*

$$b_j(n) = \left\lfloor \frac{n}{2^j} \right\rfloor - 2 \left\lfloor \frac{n}{2^{j+1}} \right\rfloor. \quad (3.1)$$

Proof. Write $\lfloor n/2^j \rfloor = 2q + r$ with $r \in \{0, 1\}$. Then $q = \lfloor n/2^{j+1} \rfloor$, while r is exactly the j -th binary digit. $\square$

Since $(-1)^b = 1 - 2b$ for $b \in \{0, 1\}$, the definition factorizes over coordinates.

Theorem 3.2 (Digit product and Boolean polynomial). *For every $n \geq 0$,*

$$\boxed{\varepsilon_n = \prod_{j \geq 0} (1 - 2b_j(n)), \quad \tau(n) = \frac{1 - \prod_{j \geq 0} (1 - 2b_j(n))}{2}.} \quad (3.2)$$

Only finitely many factors differ from 1. If J is any finite set containing the support of the binary expansion, then

$$\tau(n) = \sum_{\emptyset \neq S \subseteq J} (-2)^{|S|-1} \prod_{j \in S} b_j(n). \quad (3.3)$$

Proof. Multiply $(-1)^{b_j} = 1 - 2b_j$ over all positions. Expanding the product in the zero–one conversion (1.3) gives (3.3). $\square$

Substitution of (3.1) gives a digit-free, floor-only product:

$$\varepsilon_n = \prod_{j \geq 0} \left(1 - 2 \left\lfloor \frac{n}{2^j} \right\rfloor + 4 \left\lfloor \frac{n}{2^{j+1}} \right\rfloor \right). \quad (3.4)$$

Every factor is $1 - 2b_j(n) \in \{1, -1\}$.

Formal versions. The exact digit extraction (3.1), with the bit indicator on the left, is `Fabius.testBit_toNat_eq_div_sub_two_mul_div`; its intermediate difference-equals-remainder form is `Fabius.div_two_pow_sub_two_mul_div_two_pow_succ`. The exhaustive four-theorem API of `BinaryDigitFloor.lean` also exposes the two helpers `Fabius.div_two_pow_succ_eq_div_div` and `Fabius.sub_two_mul_div_two`. The corpus carried the equivalent *residue* form first – `Fabius.testBit_iff_div_two_pow_mod_two` in `DigitCharacterCongruence.lean`, the j -th bit is set exactly when $\lfloor n/2^j \rfloor$ is odd – and the difference form was added so that this display could be cited rather than paraphrased.

3.2 Dyadic floor and fractional-part sums

Lemma 3.3 (Legendre floor sum). *For every $n \geq 0$,*

$$\sum_{j \geq 1} \left\lfloor \frac{n}{2^j} \right\rfloor = n - \text{wt}_2(n), \quad (3.5)$$

$$\sum_{j \geq 0} \left\lfloor \frac{n}{2^j} \right\rfloor = 2n - \text{wt}_2(n). \quad (3.6)$$

Proof. Insert the binary expansion and interchange finite sums:

$$\begin{aligned} \sum_{j \geq 1} \left\lfloor \frac{n}{2^j} \right\rfloor &= \sum_{j \geq 1} \sum_{\ell \geq j} b_\ell(n) 2^{\ell-j} \\ &= \sum_{\ell \geq 1} b_\ell(n) (2^\ell - 1) = n - \sum_{\ell \geq 0} b_\ell(n). \end{aligned}$$

Adding the $j = 0$ term gives the second identity. $\square$

Theorem 3.4 (Floor-sum formulae). *For every $n \geq 0$,*

$$\varepsilon_n = (-1)^{\sum_{j \geq 0} \lfloor n/2^j \rfloor}, \quad \tau(n) \equiv \sum_{j \geq 0} \left\lfloor \frac{n}{2^j} \right\rfloor \pmod{2}. \quad (3.7)$$

Equivalently,

$$\varepsilon_n = \prod_{j \geq 0} (-1)^{\lfloor n/2^j \rfloor}. \quad (3.8)$$

Proof. By (3.6), the exponent differs from $\text{wt}_2(n)$ by an even integer. $\square$

Formal versions. `ThueMorseDigits.lean` proves the additive Legendre formula $\text{wt}_2(n) + \nu_2(n!) = n$ (`Fabius.binaryWeight_add_padicValNat_factorial`), and the floor sums as finite identities for every cutoff strictly above $\log_2 n$ (`Fabius.binaryWeight_add_sum_div_two_pow`, `Fabius.sum_div_two_pow_add_binaryWeight`), the floor-parity formula (3.7) under the same cutoff hypothesis (`Fabius.thueMorseSign_eq_neg_one_pow_sum_div`), the bound $\text{wt}_2(n) \leq n$ (`Fabius.binaryWeight_le_self`), and the bit-extraction bridge $\text{wt}_2(n) = \sum_{j < m} b_j(n)$ (`Fabius.binaryWeight_eq_sum_testBit`) together with the digit product (3.2) (`Fabius.thueMorseSign_eq_prod_testBit`). Both finite statements require $n < 2^m$, so the window contains every nonzero bit. Natural-number subtraction being awkward, every Lean statement is phrased additively; the prose forms above are recovered by moving one term across.

There is also a genuinely infinite real-valued representation whose sum is nevertheless always an integer.

Proposition 3.5 (Fractional-part digit sum). *For every $n \geq 0$,*

$$\boxed{\text{wt}_2(n) = \sum_{j \geq 1} \left\{ \frac{n}{2^j} \right\}}. \quad (3.9)$$

Consequently

$$\varepsilon_n = \cos \left(\pi \sum_{j \geq 1} \left\{ \frac{n}{2^j} \right\} \right), \quad \tau(n) = \frac{1 - \varepsilon_n}{2}. \quad (3.10)$$

Formal version. Equation (3.9) is `Fabius.tsum_fract_natCast_div_two_pow` in `DigitDefectCounting.lean`, $\sum_{h \geq 0} \{N/2^h\} = S_2(N)$; the index $h = 0$ contributes nothing, so the sum there and the sum here agree, and S_2 is wt_2 by definition of the latter. It sits in a defect-counting file rather than a Thue–Morse one because that is the role it plays: it is the regression test that fixes the sign of the weighted defect series `Fabius.tsum_mul_natCast_div_two_pow` against Legendre’s classical count. The pointwise ingredient, $\{N/2^h\} = (N \bmod 2^h)/2^h$, is `Fabius.fract_natCast_div_two_pow`.

Proof. The geometric series gives $\sum_{j \geq 1} n/2^j = n$. Subtract (3.5) termwise:

$$\sum_{j \geq 1} \left\{ \frac{n}{2^j} \right\} = n - \sum_{j \geq 1} \left\lfloor \frac{n}{2^j} \right\rfloor = \text{wt}_2(n).$$

□

3.3 Square waves and a sine-sign product

For nonintegral x ,

$$\text{sgn}(\sin \pi x) = (-1)^{\lfloor x \rfloor}. \quad (3.11)$$

The half-unit shift avoids every zero.

Theorem 3.6 (Rademacher–sine formula). *For every $n \geq 0$,*

$$\boxed{\varepsilon_n = \prod_{j \geq 0} \text{sgn} \left(\sin \frac{\pi(n + \frac{1}{2})}{2^j} \right)}. \quad (3.12)$$

All sufficiently large factors are +1. Moreover the j -th factor is exactly $(-1)^{b_j(n)}$.

Formal version. The infinite identity (3.12) is `Fabius.thueMorseSign_eq_tprod_sign_sin`; its finite-window form is `Fabius.thueMorseSign_eq_prod_sign_sin`, and the exact factor formula is `Fabius.sign_sin_rademacherPoint`, all in `RademacherSine.lean`. The same module proves the unconditional factored square wave `Fabius.sin_pi_mul_eq_neg_one_zpow_floor`, its off-integer sign form `Fabius.sign_sin_pi_mul`, positivity of the fractional factor `Fabius.sin_pi_mul_fract_pos`, the floor and nonintegrality facts `Fabius.floor_rademacherPoint` and `Fabius.fract_rademacherPoint_ne_zero`, and eventual triviality `Fabius.sign_sin_rademacherPoint_eq_one_of_lt`. Thus the displayed infinite product has finite multiplicative support and needs no analytic convergence hypothesis.

Proof. The number $(n + \frac{1}{2})/2^j$ is never integral and has the same floor as $n/2^j$. Apply (3.11), then observe from (3.1) that $(-1)^{\lfloor n/2^j \rfloor} = (-1)^{b_j(n)}$. The factors are positive once $2^j > n + \frac{1}{2}$. $\square$

Lean crosswalk (exhaustive). The nine public theorems of `RademacherSine.lean` now close this square-wave and sine-sign layer exactly. The square-wave group is `Fabius.sin_pi_mul_eq_neg_one_zpow_floor`, the unconditional factored identity valid even at integers; `Fabius.sin_pi_mul_fract_pos`, whose precise hypothesis is $\text{fract}(x) \neq 0$; and `Fabius.sign_sin_pi_mul`, with that same nonzero-fractional-part hypothesis. The half-shifted factor group is `Fabius.floor_rademacherPoint`, `Fabius.fract_rademacherPoint_ne_zero`, `Fabius.sign_sin_rademacherPoint`, and `Fabius.sign_sin_rademacherPoint_eq_one_of_lt`; the first three hold for every $n, j \in \mathbb{N}_0$, while the last assumes exactly $n < 2^j$ and proves that the factor is $+1$. Finally, `Fabius.thueMorseSign_eq_prod_sign_sin` is the finite product under exactly $n < 2^m$, and `Fabius.thueMorseSign_eq_tprod_sign_sin` is the displayed total infinite-product statement for every n . Its proof is finite-support bookkeeping—the preceding eventual-one theorem reduces $\prod'$ to a finite product—so this result asserts neither convergence of a genuinely infinite nontrivial product nor any unshifted sine formula.

3.4 The sign of the Fabius sinc product

Define $\text{sinc}_{\text{rad}} y = \sin y/y$, with $\text{sinc}_{\text{rad}} 0 = 1$, and

$$\Phi(x) = \prod_{j=0}^{\infty} \text{sinc}_{\text{rad}}\left(\frac{\pi x}{2^j}\right). \quad (3.13)$$

The product converges locally uniformly because the logarithm of the tail is controlled by $\sum_j x^2/4^j$. It is, up to normalization conventions, the infinite Fourier product naturally associated with the Rvachev up-function.

Corollary 3.7 (Interval sign of the sinc product). *If $N \in \mathbb{N}_0$ and $N < x < N + 1$, then*

$$\text{sgn } \Phi(x) = \varepsilon_N. \quad (3.14)$$

In particular,

$$\varepsilon_n = \text{sgn } \Phi\left(n + \frac{1}{2}\right). \quad (3.15)$$

Proof. For $x > 0$, every sinc denominator is positive. No multiple of 2^j lies strictly between the consecutive integers N and $N + 1$, so $\lfloor x/2^j \rfloor = \lfloor N/2^j \rfloor$. The numerator signs therefore multiply to $(-1)^{\sum_j \lfloor N/2^j \rfloor} = \varepsilon_N$ by theorem 3.4. $\square$

Formal versions. The product (3.13) is `Fabius.rvachevFourierProduct` (`Basic.lean`). The interval sign (3.14) is now the exact identity

$$\Phi(x) = \varepsilon_N |\Phi(x)| \quad (N < x < N + 1),$$

proved by `Fabius.rvachevFourierProduct_eq_thueMorse_sign_mul_norm` in `ThueMorseLobeSign.lean`. In particular, positivity of the norm off the integer lattice gives the displayed sign and the half-integer sample without an extra analytic assumption.

The same sinc product now has a complete formal zero-set statement:

$$\Phi(z) = 0 \iff \exists m \in \mathbb{Z} \setminus \{0\}, z = m,$$

namely `Fabius.rvachevFourierProduct_eq_zero_iff` in `FourierProduct.lean`. Its canonical product is

$$\Phi(z) = \prod_{r=0}^{\infty} \left(1 - \frac{z^2}{(r+1)^2}\right)^{1+\nu_2(r+1)},$$

formalized as `Fabius.rvachevFourierProduct_eq_canonical` in `SincCanonicalProduct.lean`. For each positive integer m , the older real punctured limit

$$\frac{|\Phi(x)|}{|x-m|^{1+\nu_2(m)}} \longrightarrow C_m > 0 \quad (x \rightarrow m, x \neq m)$$

is `Fabius.tendsto_norm_div_pow_zero_order` in `PhiZeroOrder.lean`. The complex theory is now stronger. For every nonzero $m \in \mathbb{Z}$, put

$$d = 1 + \nu_2(|m|), \quad q = \text{divMaxPow}(|m|, 2), \quad \sigma = \text{sign}(m).$$

Then, for every $w \in \mathbb{C}$,

$$(m+w)^d \Phi(m+w) = -w^d \left(\prod_{h < d} \text{sinc}_{\text{rad}}\left(\frac{\pi w}{2^h}\right) \right) \Phi\left(\frac{q}{2} + \sigma \frac{w}{2^d}\right).$$

This total denominator-cleared identity is `Fabius.rvachevFourierProduct_int_add_factorization` in `IntegerZeroLocalFactorization.lean`. Dividing on a neighborhood of zero produces an analytic nonvanishing cofactor U_m , and Lean proves

$$\text{ord}_{z=m} \Phi = d, \quad \Phi^{(j)}(m) = 0 \ (j < d), \quad \Phi^{(d)}(m) = -\frac{d!}{m^d} \Phi(q/2) \neq 0.$$

More generally, every higher jet is retained without factorial division:

$$\Phi^{(d+r)}(m) = \binom{d+r}{d} d! U_m^{(r)}(0) \quad (r \in \mathbb{N}_0).$$

These are `Fabius.analyticOrderAt_rvachevFourierProduct_int`, `Fabius.iteratedDeriv_rvachevFourierProduct_int_eq_zero_of_lt`, `Fabius.iteratedDeriv_rvachevFourierProduct_int_add_order`, `Fabius.iteratedDeriv_rvachevFourierProduct_int`, and `Fabius.iteratedDeriv_rvachevFourierProduct_int_ne_zero` in `IntegerZeroAnalyticOrder.lean`. That module also exports the centered germ and cofactor analyticity away from its totalized denominator point. The reusable engine in `LeadingJet.lean` is vector-valued over every nontrivially normed field and computes every index k , not just the first surviving derivative; it uses the division-free coefficient $\binom{k}{n} n!$ for a factor of order n . On the real axis the origin is the unique unit-modulus point:

$$|\Phi(x)| = 1 \iff x = 0, \quad x \neq 0 \implies |\Phi(x)| < 1.$$

These are `Fabius.norm_rvachevFourierProduct_eq_one_iff` and `Fabius.norm_rvachevFourierProduct_lt_one_of_ne_zero`. For every $m \geq 1$, the modulus also has exactly one maximizer in each positive side lobe $(m, m+1)$ and reflected negative side lobe $(-(m+1), -m)$,

by `Fabius.existsUnique_isMaxOn_lobe` and `Fabius.existsUnique_isMaxOn_neg_lobe` in `RvachevProductContinuity.lean`.

The product-level renormalization is available independently of the Fourier-transform wrapper:

$$\Phi(z) = \text{sinc}_{\text{rad}}(\pi z) \Phi(z/2), \quad \Phi(z) = \left(\prod_{j < n} \text{sinc}_{\text{rad}}(\pi z/2^j) \right) \Phi(z/2^n).$$

These are `Fabius.rvachevFourierProduct_scaling` and `Fabius.rvachevFourierProduct_shell` in `RenormalizationIdentity.lean`. Beyond these structural identities, `SincProductShells.lean` proves the exact dyadic shell factorization of the modulus: for real $y \neq 0$ and every $k \geq 0$,

$$|\Phi(2^k y)| = \frac{\prod_{j < k} |\sin(2^{j+1} \pi y)|}{2^{k(k+1)/2} (\pi |y|)^k} |\Phi(y)| \quad (3.16)$$

(`Fabius.norm_rvachevFourierProduct_two_pow_mul`): dilating the argument by 2^k creates exactly the k sine factors of the corresponding dyadic shell. Clearing the removable denominator gives the stronger total form

$$2^{k(k+1)/2} (\pi |y|)^k |\Phi(2^k y)| = \left(\prod_{j < k} |\sin(2^{j+1} \pi y)| \right) |\Phi(y)|, \quad (3.17)$$

valid for every real y , including $y = 0$, and also for the empty shell $k = 0$. This is `Fabius.norm_rvachevFourierProduct_two_pow_mul_cross`; its only zero-sensitive input is the total identity `Fabius.abs_mul_norm_complexSinc_ofReal`.

Bounding the lacunary sine numerator by the formal Gelfond bound `Fabius.abs_prod_sin_two_pow_le` (`GelfondLogisticBound.lean`) — the sine-product form of (15.8), with the optimal rate $(\sqrt{3}/2)^k$ and the slightly larger constant $\sqrt{5/3}$ in place of $2/\sqrt{3}$ — turns (3.16) into the decay bound

$$|\Phi(2^k y)| \leq \sqrt{\frac{5}{3}} \frac{(\sqrt{3}/2)^k}{2^{k(k+1)/2} (\pi |y|)^k} |\Phi(y)| \quad (3.18)$$

(`Fabius.norm_rvachevFourierProduct_two_pow_mul_le`), valid along every dyadic ray at once, with no hypothesis on y beyond $y \neq 0$. The triangular power $2^{-k(k+1)/2}$ left standing in the denominator is the universal $\exp(-(\log x)^2/(2 \log 2))$ factor of the Fabius–Rvachev decay theory; the accompanying geometric factor is what the companion reports convert into the spectrum of decay exponents.

The facade `RvachevFourierDecay.lean` identifies this standalone product with the actual Fourier transform of every bounded Fabius solution by `Fabius.rvachevFourier_eq_rvachevFourierProduct`. In particular, (3.17) transfers verbatim as `Fabius.norm_rvachevFourier_two_pow_mul_cross`. The same facade records the strict global peak

$$|\Phi(x)| = 1 \iff x = 0, \quad (3.19)$$

with strict inequality away from zero, as `Fabius.norm_rvachevFourier_eq_one_iff` and `Fabius.norm_rvachevFourier_lt_one_of_ne_zero` for the actual transform.

There is also an exact L^2 identity for the logarithm of the normalized dyadic shell. For every $n \geq 0$, including the empty product,

$$\int_0^1 \log^2 \left| \prod_{j < n} 2 \sin(\pi 2^j t) \right| dt = \frac{\pi^2}{4} n - \frac{\pi^2}{3} (1 - 2^{-n}). \quad (3.20)$$

The product form is `Fabius.integral_sq_log_prod_two_sin_all` in `BirkhoffProductBridge.lean`; the doubling-cocycle identity and its finite covariance bookkeeping are `Fabius.integral_sq_birkhoff_sum_all` and `Fabius.variance_closed_form_all` in `BirkhoffVariance.lean` and `VarianceBookkeeping.lean`.

The sine-product bound is in turn a corollary of a statement in which no trigonometry appears at all. Along *any* orbit of the full logistic map $v \mapsto 4v(1 - v)$ whose *starting point* lies in $[0, 1]$ – a hypothesis on u_0 only, the range at every later index being supplied by invariance of $[0, 1]$ – the running product satisfies

$$\prod_{j < n} u_j \leq \frac{5}{3} \left(\frac{3}{4}\right)^n \quad (3.21)$$

(`Fabius.logistic_prod_le_of_init`, `GelfondLogisticBound.lean`), by telescoping the potential $V(v) = 1 + 2v/3$ through the exact inequality $3V(v) - 4vV(4v(1 - v)) = (2v + 1)(4v - 3)^2/3 \geq 0$; the constant $5/3$ is the bound $V(u_0) \leq V(1) = 5/3$, so it too comes from $u_0 \leq 1$ alone. Squared sines at doubling angles form such an orbit (`Fabius.sin_sq_two_mul`), which is the only input the Gelfond bound needs. The geometric rate in (3.21) is exact already at this abstract level: $3/4$ is the interior fixed point of the map, and the constant orbit it generates has running product exactly $(3/4)^n$ (`Fabius.logistic_prod_eq_of_init_three_quarters`), so no smaller geometric rate can hold along all orbits whatever the constant in front. On the trigonometric side the same fixed point is attained by every seed with $\sin^2(\pi t) = 3/4$ (`Fabius.prod_sin_sq_two_pow_eq_of_sin_sq_eq`), the Gelfond angle $t = 1/3$ among them (`Fabius.prod_sin_sq_two_pow_third`).

3.5 The ruler transition

Incrementing n flips all trailing one-bits to zero and the next zero-bit to one. The number of erased one-bits is $\nu_2(n + 1)$.

Theorem 3.8 (Successor and ruler formulae). *For every $n \geq 0$,*

$$\text{wt}_2(n + 1) - \text{wt}_2(n) = 1 - \nu_2(n + 1), \quad (3.22)$$

$$\varepsilon_{n+1} = (-1)^{1+\nu_2(n+1)} \varepsilon_n, \quad (3.23)$$

$$(-1)^{\nu_2(n+1)} = -\varepsilon_n \varepsilon_{n+1}. \quad (3.24)$$

Consequently

$$\varepsilon_n = \prod_{k=1}^n (-1)^{1+\nu_2(k)}. \quad (3.25)$$

Proof. The bit transition changes the weight by one minus the number of trailing ones. Exponentiate (-1) , then telescope from $\varepsilon_0 = 1$. $\square$

Formal versions. The successor law in additive form $\text{wt}_2(n + 1) + \nu_2(n + 1) = \text{wt}_2(n) + 1$ is `Fabius.binaryWeight_succ_add_padicValNat`; the adjacent-sign readout (3.24) is `Fabius.thueMorseSign_mul_succ`, the recursion (3.23) is `Fabius.thueMorseSign_succ_eq`, and the telescoped ruler product (3.25) is `Fabius.thueMorseSign_eq_prod_ruler`, all in `ThueMorseValuation.lean`.

4 Factorial valuations, central binomials, Catalan numbers, and carries

4.1 Legendre’s formula

The floor sum in [lemma 3.3](#) counts powers of two in a factorial.

Theorem 4.1 (Factorial valuation). *For every $n \geq 0$,*

$$\nu_2(n!) = \sum_{j \geq 1} \left\lfloor \frac{n}{2^j} \right\rfloor = n - \text{wt}_2(n). \quad (4.1)$$

Hence

$$\varepsilon_n = (-1)^{n-\nu_2(n!)} = (-1)^{n+\nu_2(n!)}, \quad \tau(n) \equiv n + \nu_2(n!) \pmod{2}. \quad (4.2)$$

Proof. Every multiple of 2 contributes at least one factor, every multiple of 4 contributes one additional factor, and so on. This gives the floor sum; lemma 3.3 gives the digit form. Addition and subtraction agree modulo two. $\square$

The telescoping product (3.25) returns to this formula because $\sum_{k=1}^n \nu_2(k) = \nu_2(n!)$.

Formal versions. Legendre’s formula (4.1) is `Fabius.binaryWeight_add_padicValNat_factorial` in `ThueMorseDigits.lean`, stated additively as $\text{wt}_2(n) + \nu_2(n!) = n$ so that no subtraction in $\mathbb{N}_0$ is needed. The floor form (3.5) is `Fabius.binaryWeight_add_sum_div_two_pow` and the all-index variant (3.6) is `Fabius.sum_div_two_pow_add_binaryWeight`, both in the same file. Each floor statement carries one hypothesis the display leaves implicit: the sum is truncated at any $b > \log_2 n$, which is where its tail is already zero.

4.2 A single central binomial coefficient

Theorem 4.2 (Central-binomial valuation). *For every $n \geq 0$,*

$$\nu_2 \binom{2n}{n} = \text{wt}_2(n). \quad (4.3)$$

Consequently

$$\varepsilon_n = (-1)^{\nu_2 \binom{2n}{n}}, \quad \tau(n) = \nu_2 \binom{2n}{n} \pmod{2}. \quad (4.4)$$

Equivalently,

$$\tau(n) = \sin^2 \left(\frac{\pi}{2} \nu_2 \binom{2n}{n} \right), \quad \varepsilon_n = \cos \left(\pi \nu_2 \binom{2n}{n} \right). \quad (4.5)$$

Proof. Use (4.1) and $\text{wt}_2(2n) = \text{wt}_2(n)$:

$$\begin{aligned} \nu_2 \binom{2n}{n} &= \nu_2((2n)!) - 2\nu_2(n!) \\ &= (2n - \text{wt}_2(2n)) - 2(n - \text{wt}_2(n)) = \text{wt}_2(n). \end{aligned}$$

$\square$

Formal versions. The exhaustive three-theorem public API of `CentralBinomialValuation.1` can be `Fabius.padicValNat_two_centralBinom`, `Fabius.thueMorseSign_eq_neg_one_pow_centralBinom`, and `Fabius.padicValNat_two_centralBinom_eq_zero_iff`. The last theorem states that the valuation vanishes exactly when the binary weight is zero, which occurs only at $n = 0$; a positive power of two has valuation one, not zero. The first identity is the $a = b = n$ case of the additive Kummer theorem `Fabius.binaryWeight_add_addChoose_padicValNat`, with $\text{wt}_2(2n) = \text{wt}_2(n)$ cancelling one copy of the weight. Thus the Lean proof does not depend on a separate factorial-valuation argument. The $\tau(n)$ -mod-two and trigonometric displays are immediate report-level corollaries, not separate named declarations in this module.

The shortest Pascal replacement. The source formula inspects all $n + 1$ entries of row n and takes an exact logarithm. Formula (4.4) inspects the two-adic order of one entry in row $2n$, and in fact recovers the entire digit sum rather than only its parity.

4.3 Catalan valuations

Let

$$\text{Cat}_n = \frac{1}{n+1} \binom{2n}{n}. \quad (4.6)$$

Corollary 4.3 (Catalan formula). *For every $n \geq 0$,*

$$\nu_2(\text{Cat}_n) = \text{wt}_2(n+1) - 1, \quad (4.7)$$

and hence

$$\boxed{\varepsilon_{n+1} = -(-1)^{\nu_2(\text{Cat}_n)}}. \quad (4.8)$$

Proof. By [theorem 4.2](#), $\nu_2(\text{Cat}_n) = \text{wt}_2(n) - \nu_2(n+1)$. Equation (3.22) rewrites this as $\text{wt}_2(n+1) - 1$. $\square$

Formal versions. `Fabius.catalan_padicValNat_two_add_one` (stated additively as $\nu_2(\text{Cat}_n) + 1 = \text{wt}_2(n+1)$) and the sign form `Fabius.thueMorseSign_succ_eq_neg_catalan`, in `ThueMorseValuation.lean`. The Lean proof derives the positivity of Catalan numbers inline from $(n+1)\text{Cat}_n = \binom{2n}{n}$ rather than importing it.

4.4 Kummer’s carry cocycle

Theorem 4.4 (Addition with carries). *For all $a, b \geq 0$,*

$$\nu_2 \binom{a+b}{a} = \text{wt}_2(a) + \text{wt}_2(b) - \text{wt}_2(a+b). \quad (4.9)$$

This is the number of carry operations in the binary addition of a and b , including propagated carries. Therefore

$$\boxed{\varepsilon_{a+b} = \varepsilon_a \varepsilon_b (-1)^{\nu_2 \binom{a+b}{a}}}. \quad (4.10)$$

Equivalently,

$$\tau(a+b) \equiv \tau(a) + \tau(b) + \nu_2 \binom{a+b}{a} \pmod{2}. \quad (4.11)$$

Proof. Legendre’s formula gives

$$\nu_2 \binom{a+b}{a} = (a+b - \text{wt}_2(a+b)) - (a - \text{wt}_2(a)) - (b - \text{wt}_2(b)),$$

which is (4.9). Rearranging modulo two and raising (-1) to both sides gives the sign formula. The carry interpretation is the binary case of Kummer’s theorem. $\square$

The no-carry case $a \& b = 0$ makes the valuation zero and recovers ordinary multiplicativity. In particular, [theorem 2.2](#) is the separated-support case.

Corollary 4.5 (Multinomial carry formula). *For $a_1, \dots, a_r \geq 0$ and $N = a_1 + \dots + a_r$,*

$$\nu_2 \binom{N}{a_1, \dots, a_r} = \sum_{j=1}^r \text{wt}_2(a_j) - \text{wt}_2(N), \quad (4.12)$$

$$\varepsilon_N = \left(\prod_{j=1}^r \varepsilon_{a_j} \right) (-1)^{\nu_2 \binom{N}{a_1, \dots, a_r}}. \quad (4.13)$$

Proof. Apply Legendre’s formula to $N!/(a_1! \dots a_r!)$ and then take parity. $\square$

Formal versions. Kummer’s identity (4.9) is `Fabius.binaryWeight_add_addChoose_padicValNat` in `ThueMorseValuation.lean`, again stated additively as $\text{wt}_2(a + b) + \nu_2\binom{a+b}{a} = \text{wt}_2(a) + \text{wt}_2(b)$; the sign cocycle (4.10) is `Fabius.thueMorseSign_add_valuation` in `TwoAdic.lean`, and the borrow form for subtraction is `Fabius.subChoose_padicValNat_two_add`. The multinomial valuation (4.12) is formalized at an *arbitrary* prime rather than only at two, as `Fabius.sum_digitSum_eq_sub_one_mul_padicValNat_multinomial` in `ThueMorseMultinomial.lean`, whose $p = 2$ reading is the display; its sign form (4.13) follows by taking parity but is not separately stated there.

4.5 Subtraction and borrow counts

Kummer’s theorem has a mirror statement for subtraction. For $0 \leq b \leq a$, the same Legendre computation applied to $\binom{a}{b} = a!/(b!(a-b)!)$ gives

$$\nu_2\binom{a}{b} = \text{wt}_2(b) + \text{wt}_2(a - b) - \text{wt}_2(a), \quad (4.14)$$

which counts the borrow operations, including propagated borrows, in the binary subtraction $a - b$. Therefore

$$\boxed{\varepsilon_{a-b} = \varepsilon_a \varepsilon_b (-1)^{\nu_2\binom{a}{b}}.} \quad (4.15)$$

Addition and subtraction thus carry the same sign correction, read from the two-adic valuation of the binomial coefficient whose top entry is the larger of the numbers involved; (4.10) is recovered from (4.15) by renaming $a + b$ as the top entry.

Formal versions. Kummer’s carry formula and cocycle were already in the corpus (`Fabius.addChoose_padicValNat_two`, `Fabius.thueMorseSign_add_valuation`, in `TwoAdic.lean`); `ThueMorseValuation.lean` adds the truncation-free additive form `Fabius.binaryWeight_add_addChoose_padicValNat`, the weight subadditivity $\text{wt}_2(a + b) \leq \text{wt}_2(a) + \text{wt}_2(b)$ (`Fabius.binaryWeight_add_le`, itself proved by a four-way parity induction), and the borrow pair `Fabius.subChoose_padicValNat_two_add`, `Fabius.thueMorseSign_sub_valuation`.

5 Pascal support, residue formulae, multinomials, and Boolean inversion

5.1 Lucas’s theorem as a submask criterion

Theorem 5.1 (Lucas parity criterion). *For $0 \leq k \leq n$,*

$$\binom{n}{k} \equiv 1 \pmod{2} \iff k \preceq_2 n. \quad (5.1)$$

Equivalently,

$$\mathbf{1}_{k \preceq_2 n} = \frac{1 - (-1)^{\binom{n}{k}}}{2}. \quad (5.2)$$

Proof. In $\mathbb{F}_2[z]$, Frobenius gives

$$(1 + z)^n = \prod_{j: b_j(n)=1} (1 + z)^{2^j} = \prod_{j: b_j(n)=1} (1 + z^{2^j}).$$

The exponents occurring in the expansion are exactly the submasks of n , and their coefficients are one modulo two. $\square$

Define the odd-support polynomial of row n by

$$Q_n(z) = \sum_{k=0}^n \mathbf{1}_{\binom{n}{k} \text{ odd}} z^k. \quad (5.3)$$

The proof gives the complete factorization

$$Q_n(z) = \prod_{j: b_j(n)=1} (1 + z^{2^j}). \quad (5.4)$$

Corollary 5.2 (Glaisher–Fine count). *The number $\nu(n)$ of odd entries in row n is*

$$\nu(n) = 2^{\text{wt}_2(n)}. \quad (5.5)$$

Proof. Set $z = 1$ in (5.4). □

Formal versions. `ThueMorseLucasSupport.lean` proves [theorem 5.1](#) in three forms: `Fabius.odd_choose_iff_testBit` (every set bit of k is set in n), `Fabius.odd_choose_iff_land` ($k \& n = k$), and `Fabius.odd_choose_iff_bitSupport_subset`, by strong induction through the parity recursions of Pascal rows rather than through $\mathbb{F}_2[z]$. The supporting bit dictionary is centralized in `ThueMorseBitSupport.lean`: a sum of distinct two-powers has exactly the prescribed binary support by `Fabius.testBit_sum_two_pow`, stated for an arbitrary finite set of bit positions, together with the two inverse reconstruction laws. The Glaisher–Fine count (5.5) has long been formal as `Fabius.card_oddBinomialIndices` (`Parity.lean`); the odd row positions are identified with the submask cube by `Fabius.oddBinomialIndices_eq_image_powerset`.

The source formula is now transparent. Since even entries contribute $+1$ and odd entries contribute -1 ,

$$\mathcal{X}(n) = n + 1 - 2\nu(n) = n + 1 - 2^{\text{wt}_2(n)+1}, \quad (5.6)$$

which is exactly (1.5).

5.2 A logarithm-free Pascal-row formula modulo three

Let $\text{rem}_3(x) \in \{0, 1, 2\}$ denote the least nonnegative residue modulo three. Powers of two alternate between one and two modulo three.

Theorem 5.3 (Odd-count residue formula). *For every $n \geq 0$,*

$$\tau(n) = \text{rem}_3(\nu(n)) - 1, \quad \varepsilon_n = 3 - 2\text{rem}_3(\nu(n)). \quad (5.7)$$

The residue is always 1 or 2, never 0.

Proof. By (5.5), $\nu(n) = 2^{\text{wt}_2(n)}$. Its residue is one when $\text{wt}_2(n)$ is even and two when it is odd. □

Formal versions. Both displays of (5.7) are in `ThueMorseBinomialLog.lean`, with $\nu(n)$ read as the cardinality of `Fabius.oddBinomialIndices`: the bit form $\tau(n) = \text{rem}_3(\nu(n)) - 1$ is `Fabius.thueMorseBit_eq_card_oddBinomialIndices_mod_three_sub_one`, its rearrangement $\text{rem}_3(\nu(n)) = \tau(n) + 1$ is `Fabius.card_oddBinomialIndices_mod_three_eq_bit_add_one`, and the sign form $\varepsilon_n = 3 - 2\text{rem}_3(\nu(n))$ is `Fabius.thueMorseSign_eq_three_sub_two_mul_card_oddBinomialIndices_mod_three`. The closing remark that the residue is never zero is `Fabius.card_oddBinomialIndices_mod_three_ne_zero`. The arithmetic underneath all four is the two-cycle of powers of two modulo three, `Fabius.two_pow_mod_three` (in `SignedRowSum.lean`; `ThueMorseBinomialLog.lean` keeps a private copy), $2^r \bmod 3 = (r \bmod 2) + 1$.

A form using exactly the row signs of the source formula can be written without the odd-count notation. Put

$$S(n) = \sum_{k=1}^n (-1)^{\binom{n}{k}}, \quad S(0) = 0. \quad (5.8)$$

Since the omitted $k = 0$ term is -1 ,

$$S(n) = n + 2 - 2^{\text{wt}_2(n)+1}. \quad (5.9)$$

Corollary 5.4 (Binomial residue and sine formulae). *For every $n \geq 0$,*

$$\tau(n) = \text{rem}_3(2n + S(n)), \quad (5.10)$$

$$\tau(n) = \frac{4}{3} \sin^2\left(\frac{\pi}{3}(n - S(n))\right). \quad (5.11)$$

Proof. Equation (5.9) gives $2n + S(n) = 3n + 2 - 2^{\text{wt}_2(n)+1}$, whose residue is zero or one according as the weight is even or odd. Also $n - S(n) = 2^{\text{wt}_2(n)+1} - 2$, congruent to zero or two modulo three; the squared sine is therefore zero or $3/4$. $\square$

Formal versions. The signed row sum is `Fabius.signedRowSum` in `SignedRowSum.lean`, and (5.9) is `Fabius.signedRowSum_eq`, obtained from the full-row sum $\sum_{k=0}^n (-1)^{\binom{n}{k}} = (n+1) - 2^{\text{wt}_2(n)+1}$ (`Fabius.sum_range_neg_one_pow_choose`) by removing the $k = 0$ term. The residue display (5.10) is `Fabius.thueMorseBit_eq_two_mul_add_signedRowSum_emod_three`, and the sine display (5.11) is `Fabius.thueMorseBit_eq_sin_sq_pi_sub_signedRowSum_div_three`; the latter rests on $n - S(n) = 6k + 2\tau(n)$ with k exhibited explicitly (`Fabius.two_pow_succ_sub_two_eq`) and on $\sin^2(2\pi/3) = 3/4$ (`Fabius.sin_sq_two_pi_div_three`).

The logarithmic and residue formulas are two shadows of the same invariant $\nu(n)$: the first reads the exponent of a power of two; the second reads that exponent's parity from the two-cycle of powers modulo three.

5.3 A one-parameter row statistic

For a scalar u , define

$$R_u(n) = \sum_{k=0}^n u^{\binom{n}{k}} \bmod 2. \quad (5.12)$$

Exactly $\nu(n)$ exponents are one and the others are zero, so

$$R_u(n) = n + 1 + (u - 1)2^{\text{wt}_2(n)}. \quad (5.13)$$

Thus $R_{-1}(n) = \mathcal{X}(n)$, while $R_0(n)$ counts the even entries. For every fixed $u \neq 1$,

$$2^{\text{wt}_2(n)} = \frac{R_u(n) - (n + 1)}{u - 1}. \quad (5.14)$$

The signed choice $u = -1$ is elegant, but structurally it is one member of a whole exact family.

5.4 Sparse Pascal probes for individual bits

Lucas's criterion at $k = 2^j$ isolates one binary coordinate:

$$\binom{n}{2^j} \equiv b_j(n) \pmod{2}. \quad (5.15)$$

Theorem 5.5 (Power-of-two column formula). *For every $n \geq 0$,*

$$\boxed{\varepsilon_n = (-1)^{\sum_{j \geq 0} \binom{n}{2^j}} = \prod_{j \geq 0} \cos\left(\pi \binom{n}{2^j}\right)}, \quad (5.16)$$

$$\tau(n) = \frac{1}{2} \left(1 - \prod_{j \geq 0} \cos\left(\pi \binom{n}{2^j}\right) \right). \quad (5.17)$$

Only the columns $2^j \leq n$ matter.

Proof. Only the parity of each binomial coefficient contributes to the exponent or cosine. Equation (5.15) turns their parity sum into $\text{wt}_2(n)$. $\square$

Formal version. The per-column probe (5.15) is Lucas’s criterion, formal in three interchangeable readings in `ThueMorseLucasSupport.lean`: the bit form `Fabius.odd_choose_iff_testBit`, the bitwise-and form `Fabius.odd_choose_iff_land` ($\binom{n}{k}$ is odd iff $k \& n = k$), and the support form `Fabius.odd_choose_iff_bitSupport_subset`. Taking $k = 2^j$ gives exactly the probe. The summed statements (5.16) and (5.17) follow by adding the parities over j , but are not separately stated; the same reduction is carried out inline in `ThueMorseIntegerLift.lean`, where the column probes are summed against Catalan parities.

Unlike the source formula, this scans only $O(\log n)$ selected entries and uses no logarithm. Its price is the distinguished set of power-of-two columns.

5.5 Odd multinomial coefficients and exact r -nomial logarithms

Fix $r \geq 2$. Let $O_r(n)$ be the number of odd coefficients in $(x_1 + \cdots + x_r)^n$.

Theorem 5.6 (Multinomial support count). *For every $n \geq 0$,*

$$\boxed{O_r(n) = r^{\text{wt}_2(n)}}. \quad (5.18)$$

Consequently

$$\boxed{\varepsilon_n = (-1)^{\log_r O_r(n)}, \quad \tau(n) = \frac{1 - (-1)^{\log_r O_r(n)}}{2}}. \quad (5.19)$$

The logarithm is exact.

Proof. In characteristic two,

$$(x_1 + \cdots + x_r)^n = \prod_{j: b_j(n)=1} (x_1^{2^j} + \cdots + x_r^{2^j}). \quad (5.20)$$

For each one-bit of n , choose independently one of the r variables to receive it. Unique binary expansions make distinct assignments produce distinct exponent vectors, so there are $r^{\text{wt}_2(n)}$ nonzero coefficients modulo two. $\square$

There are $\binom{n+r-1}{r-1}$ weak compositions of n into r parts. Define

$$\mathcal{X}_r(n) = \sum_{k_1 + \cdots + k_r = n} (-1)^{\binom{n}{k_1, \dots, k_r}}. \quad (5.21)$$

Corollary 5.7 (Signed r -nomial row formula). *For every $r \geq 2$,*

$$\binom{n+r-1}{r-1} - \mathcal{X}_r(n) = 2r^{\text{wt}_2(n)}, \quad (5.22)$$

and therefore

$$\tau(n) = \frac{1 - (-1)^{\log_r \left(\frac{\binom{n+r-1}{r-1} - \mathcal{X}_r(n)}{2} \right)}}{2}. \quad (5.23)$$

Proof. Each odd multinomial coefficient contributes -1 , each even coefficient $+1$. Subtracting the signed sum from the total count gives twice the odd support size, which is $2r^{\text{wt}_2(n)}$ by [theorem 5.6](#). $\square$

Formal versions. `ThueMorseMultinomial.lean` proves [theorem 5.6](#) as `Fabius.card_filter_odd_multinomial`, for an arbitrary finite index set, resting on `**Kummer’s criterion for multinomial coefficients**` `Fabius.odd_multinomial_iff_binaryWeight`: a multinomial coefficient is odd exactly when the binary weights of its parts are additive. The proof uses no polynomial ring: Legendre’s formula converts oddness into weight additivity, and the induction that splits off one part reduces the count to the bivariate submask enumerator at $(1, r)$. The exact logarithm (5.19) is `Fabius.thueMorseSign_eq_neg_one_pow_log_card`, and the signed row formula (5.22) is `Fabius.sum_neg_one_pow_multinomial` (with the generic count-parity bridge `Fabius.sum_neg_one_pow_eq_card_sub_two_mul`); the closed value $\binom{n+r-1}{r-1}$ of the total count is left to the reader’s `Mathlib`.

5.6 Boolean weight enumerators and Möbius inversion

The submasks of n form a Boolean lattice on its set of one-bit positions.

Theorem 5.8 (Submask weight enumerator). *For every indeterminate z ,*

$$\sum_{k \preceq_2 n} z^{\text{wt}_2(k)} = (1 + z)^{\text{wt}_2(n)}, \quad (5.24)$$

$$z^{\text{wt}_2(n)} = \sum_{k \preceq_2 n} (z - 1)^{\text{wt}_2(k)}. \quad (5.25)$$

More generally,

$$\sum_{k \preceq_2 n} x^{\text{wt}_2(k)} y^{\text{wt}_2(n-k)} = (x + y)^{\text{wt}_2(n)}. \quad (5.26)$$

Proof. For each occupied bit of n , choose whether it belongs to k , to its complement $n - k$, or is omitted in the one-variable specialization. The second formula is the first with z replaced by $z - 1$. $\square$

Important specializations are

$$\#\{k : k \preceq_2 n\} = 2^{\text{wt}_2(n)}, \quad (5.27)$$

$$\sum_{k \preceq_2 n} \varepsilon_k = \begin{cases} 1, & n = 0, \\ 0, & n > 0, \end{cases} \quad (5.28)$$

$$\varepsilon_n = \sum_{k \preceq_2 n} (-2)^{\text{wt}_2(k)}. \quad (5.29)$$

Theorem 5.9 (Boolean Möbius inversion). *If $k \preceq_2 n$, then the Möbius function of the submask poset is*

$$\mu_2(k, n) = \varepsilon_n \varepsilon_k = \varepsilon_{n \oplus k}. \quad (5.30)$$

Thus

$$g(n) = \sum_{k \preceq_2 n} f(k) \iff \boxed{f(n) = \sum_{k \preceq_2 n} \varepsilon_{n \oplus k} g(k)}. \quad (5.31)$$

Proof. The interval $[k, n]$ is the Boolean lattice on the bits of n absent from k ; its rank is $\text{wt}_2(n) - \text{wt}_2(k) = \text{wt}_2(n \oplus k)$. The usual Boolean Möbius function is $(-1)^{\text{rank}}$, and standard Möbius inversion applies. $\square$

Formal versions. `ThueMorseBooleanMöbius.lean` proves the whole lattice layer. The enumerators (5.24) and (5.26) are `Fabius.sum_oddBinomialIndices_pow_binaryWeight` and `Fabius.sum_oddBinomialIndices_pow_pow`, over any commutative semiring; (5.25) and the reconstruction (5.29) are `Fabius.pow_binaryWeight_eq_sum_oddBinomialIndices` and `Fabius.thueMorseSign_eq_sum_oddBinomialIndices_neg_two_pow`; the count (5.27) and the cancellation (5.28) were already `Fabius.card_oddBinomialIndices` and `Fabius.sum_thueMorseSign_oddBinomialIndices`. Möbius inversion itself is proved in two forms: on finsets for any additive commutative group (`Fabius.sum_powerset_mobius`, with the interval sum `Fabius.sum_filter_superset_neg_one_pow`), and numerically on submasks with the Möbius function $\varepsilon_{n \oplus k}$ (`Fabius.submask_mobius_inversion`). The canonical two-power encoding and its inverse `Fabius.bitSupport_sum_two_pow` live in `ThueMorseBitSupport.lean`; the submask complement laws `Fabius.sub_sum_two_pow` and `Fabius.xor_sum_two_pow` remain in `ThueMorseBooleanMöbius.lean`. The matrix statement (5.33) is the poset form read row by row.

Let

$$P_{n,k} = \mathbf{1}_{\binom{n}{k} \text{ odd}} = \mathbf{1}_{k \preceq_2 n}. \quad (5.32)$$

As a locally finite lower-triangular matrix,

$$\boxed{(P^{-1})_{n,k} = P_{n,k} \varepsilon_{n \oplus k}}. \quad (5.33)$$

In particular, the first column of P^{-1} is the signed Thue–Morse sequence. This is a second global Pascal encoding: ε is not only a row statistic but the Möbius inverse of the entire parity incidence matrix.

Combining (5.29), Lucas’s indicator, and the central binomial valuation gives a formula involving only ordinary binomial coefficients:

Theorem 5.10 (Nested binomial formula). *For every $n \geq 0$,*

$$\boxed{\varepsilon_n = \sum_{k=0}^n \frac{1 - (-1)^{\binom{n}{k}}}{2} (-2)^{\nu_2 \binom{2k}{k}}}. \quad (5.34)$$

Proof. The first factor is $\mathbf{1}_{k \preceq_2 n}$ by (5.2), and the exponent in the second is $\text{wt}_2(k)$ by theorem 4.2. The sum is therefore exactly (5.29). $\square$

Formal versions. (5.34) is `Fabius.thueMorseSign_eq_sum_nested_binomial` (`ThueMorseBooleanMöbius.lean`), from the reconstruction formula and the central-binomial valuation `Fabius.centralChoose_padicValNat_two`.

6 Finite and infinite generating products

6.1 The Boolean-cube mother identity

A multivariate identity contains the finite generating product, the weight enumerator, and the Walsh character formula as specializations.

Theorem 6.1 (Boolean-cube product). *For commuting indeterminates $x_0, \dots, x_{m-1}$,*

$$\sum_{n=0}^{2^m-1} \varepsilon_n \prod_{j=0}^{m-1} x_j^{b_j(n)} = \prod_{j=0}^{m-1} (1 - x_j). \quad (6.1)$$

Formal versions. The principal ingredients of the multivariable identity are formal in two pieces, and the split is the interesting part. Over any commutative semiring, `Fabius.prod_one_add_eq_sum_powerset` in `ThueMorseBooleanCube.lean` is $\prod_{j \in S} (1 + x_j) = \sum_{T \subseteq S} \prod_{j \in T} x_j$ – the display with the sum still indexed by *subsets* – and `Fabius.sum_powerset_two_pow` in `ThueMorseBitSupport.lean` is the re-indexing $T \mapsto \sum_{j \in T} 2^j$ that turns subsets of $\{0, \dots, m-1\}$ into $\{0, \dots, 2^m - 1\}$. Together with the sign and bit-support dictionary, they are the principal formal ingredients of (6.1); the arbitrary multivariable composite itself has no single named declaration. The corpus states the composite only in the one-variable specialization $x_j = z^{2^j}$, as `Fabius.prod_one_sub_pow_eq_sum_thueMorseSign`, and in the two-parameter form `Fabius.prod_one_add_mu_1_pow`, $\prod_{j < m} (1 + uz^{2^j}) = \sum_{n < 2^m} u^{\text{wt}_2(n)} z^n$, from which the finite Thue–Morse product and the weight enumerator follow by setting $u = -1$ and $z = 1$, respectively.

Proof. Identify n with its bit vector $b \in \{0, 1\}^m$. The sum factors over coordinates:

$$\sum_{b \in \{0,1\}^m} \prod_j (-x_j)^{b_j} = \prod_j (1 - x_j).$$

□

Setting all $x_j = -u$ recovers (2.18); setting $x_j = z^{2^j}$ gives the central finite polynomial.

6.2 The finite Prouhet polynomial

Define

$$P_m(z) = \sum_{n=0}^{2^m-1} \varepsilon_n z^n. \quad (6.2)$$

Theorem 6.2 (Complete finite factorization). *For every $m \geq 0$,*

$$P_m(z) = \prod_{j=0}^{m-1} (1 - z^{2^j}). \quad (6.3)$$

The empty product for $m = 0$ is 1.

Proof. Use theorem 6.1 with $x_j = z^{2^j}$. Equivalently, expanding the product chooses a subset of powers of two; its sum is an integer $n < 2^m$, uniquely, and the coefficient is $(-1)^{\#\text{chosen bits}} = \varepsilon_n$. □

Formal version. `Fabius.prod_one_sub_pow_two_pow` (`ThueMorseProductIdentity.lean`) proves (6.3) over an arbitrary commutative ring and an arbitrary ring element z : the Boolean-cube expansion (`Finset.prod_add`) is reindexed by the corpus’s dyadic-code dictionary, so each subset’s sign becomes the binary-weight sign of its code. The same module proves the

subset-sum polynomial layer of the exact-spline theory: `Fabius.one_sub_mul_prod_one_sub_pow_two_pow` shows that multiplying by the duplicated smallest factor $1 - z$ takes first differences of the sign word (the knot weights `Fabius.thueMorseWordDelta` of the finite box splines), `Fabius.one_sub_mul_prod_eq_sq_mul` records the $(1 - z)^2$ -form, and evaluation at $z = 1$ (`Fabius.sum_thueMorseSign_two_pow`) recovers the balance of every level- m sign block.

Formal version. The cosine companion of the lacunary product-to-sum identity is formal in `Lacuna ryProductToSumCosine.lean`. Shifting the affine phase by a quarter turn turns the corpus’s affine sine factorization into its cosine counterpart (`Fabius.sum_thueMorseSign_mul_cos_affine`), and hence the lacunary product with its top factor replaced by a cosine into the Thue–Morse cosine polynomial at the same odd frequencies, $2^m \cos(2^m w) \prod_{j < m} \sin(2^j w) = \sum_{n < 2^m} \varepsilon_n \cos(m\pi/2 + (2n + 1)w)$ (`Fabius.prod_sin_two_pow_mul_cos_eq_thueMorse_sum`); the rescaled z -form (`Fabius.prod_sin_inv_two_pow_mul_cos_eq_thueMorse_sum`) is the display used here. Two elementary inputs of the shell estimates are formal as well: the multiple-angle domination bounds $|\sin mx| \leq m|\sin x|$ for every m and $|\cos mx| \leq m|\cos x|$ for odd m (`Fabius.abs_sin_nat_mul_le`, `Fabius.abs_cos_nat_mul_le_of_odd`, with a sharpness witness for the parity hypothesis), and the Weierstrass product inequality $1 - \sum u_i \leq \prod (1 - u_i)$ in finite and infinite form (`Fabius.one_sub_sum_le_prod_one_sub`, `Fabius.one_sub_tsum_le_tprod_one_sub`) – neither of which was available in `Mathlib`.

Two further consequences of the same block algebra are formal. Dividing the block’s ladder factorization by $(1 - X)^{d+1}$ in $\mathbb{Z}[X]$ puts the corpus’s approximation polynomial – the q -integer product $A_d = \prod_{m=1}^d [2^m]_X$ – in closed form as a product of self-reciprocal factors with scale-decreasing multiplicity, $A_d = \prod_{i < d} (1 + X^{2^i})^{d-i}$ (`Fabius.approximationPolynomialInt_eq_ladder`, `Fabius.approximationPolynomialInt_eq_ladder’`). And the Boolean-cube master identity acquires its *bit-position* companion: where `Fabius.prod_one_add_mul_pow` grades a dyadic block by the value z^n ,

$$\prod_{j < N} (1 + y q^j) = \sum_{n < 2^N} y^{\text{wt}_2(n)} q^{\sigma(n)}, \quad \sigma(n) = \sum_{j \in \text{bitSupport}(n)} j,$$

grades it by the sum of its bit positions (`Fabius.bitPositionSum`, `Fabius.prod_one_add_mul_pow_bitPositionSum`), over any commutative semiring; the $y = 1$ and $q = 1$ faces recover the q^σ -enumerator and the binary-weight enumerator respectively. Its y^r -coefficient is now formal too: the q -binomial aggregation of a dyadic block by Hamming weight,

$$\sum_{\substack{n < 2^N \\ \text{wt}_2(n)=r}} q^{\sigma(n)} = q^{\binom{r}{2}} \begin{bmatrix} N \\ r \end{bmatrix}_q,$$

(`Fabius.sum_pow_bitPositionSum_filter_eq_gaussianBinomial`), with the $q = 1$ shadow recovering the binomial coefficient.

That shadow is now available in the generality the rest of the q -layer is stated in. `Gaussian BinomialAtOne.lean` proves $\begin{bmatrix} n \\ k \end{bmatrix}_1 = \binom{n}{k}$ over an *arbitrary* semiring – by an induction over $\mathbb{N}_0$, where the q -Pascal recurrence degenerates to Pascal’s rule, transported along the functoriality of the Gaussian coefficient (`Fabius.gaussianBinomial_nat_one`, `Fabius.gaussianBinomial_one`; the corpus previously had no $q = 1$ evaluation at all, and the graded aggregation’s own version was confined to $\mathbb{Z}$). Alongside it the q -Pochhammer symbol acquires the same specializations – $(a; 1)_n = (1 - a)^n$, and the degenerate $a = 0$, $a = 1$ cases (`Fabius.finiteQPochhammerIn_one`, `Fabius.finiteQPochhammerIn_zero_left`, `Fabius.finiteQPochhammerIn_one_left`) – and together these give the consistency check the whole layer is measured against: the finite q -binomial theorem at $q = 1$ is the ordinary binomial theorem, $\sum_{k \leq n} (-1)^k \binom{n}{k} z^k = (1 - z)^n$ (`Fabius.sum_neg_one_pow_choose_mul_pow`).

The same block algebra also inverts. The truncated product $\Theta_m = \prod_{j < m} (1 - z^{2^j})$ has the restricted binary partition counts as its convolution (Green) inverse: $\Theta_m \cdot P_m = 1$ in $\mathbb{Z}[[z]]$ (Fabius.prod_one_sub_X_two_pow_mul_binaryPartitionSeries), whence the finite convolution $\sum_r \varepsilon_r p_m(N - r) = [N = 0]$ (Fabius.thueMorseSign_binaryPartitionCount_convolution). The summation range there is $r < \min(N + 1, 2^m)$ and not merely $r < 2^m$: truncated natural subtraction sends every $r > N$ to $p_m(0) = 1$, so the naive range is false already at $N = 0$.

Finally, the exact weighted zero-counting function acquires its *defect* form. For a real weight sequence with $\sum_h |a_h| 2^{-h} < \infty$,

$$\sum_{h \geq 0} a_h \left\lfloor \frac{N}{2^h} \right\rfloor = N A(\tfrac{1}{2}) - \rho_a(N), \quad A(\tfrac{1}{2}) = \sum_h \frac{a_h}{2^h}, \quad \rho_a(N) = \sum_h a_h \left\{ \frac{N}{2^h} \right\}$$

(Fabius.tsum_mul_natCast_div_two_pow_defect, with the corpus form Fabius.sum_range_weightedScaleMultiplicity_real). Boundedness of a is not needed — the summability hypothesis alone dominates both series. At $a \equiv 1$ the identity reduces to Legendre’s $\sum_h \lfloor N/2^h \rfloor = 2N - \text{wt}_2(N)$, which is carried as a regression test (Fabius.sum_range_div_two_pow_add_digitSum, Fabius.dyadicDefect_one).

Corollary 6.3 (Coefficient and block formulae). *For every $n \geq 0$,*

$$\varepsilon_n = [z^n] \prod_{j \geq 0} (1 - z^{2^j}), \quad (6.4)$$

where only factors with $2^j \leq n$ affect the coefficient. Moreover, for $q, m \geq 0$,

$$\sum_{r=0}^{2^m-1} \varepsilon_{2^m q+r} z^r = \varepsilon_q P_m(z). \quad (6.5)$$

Proof. The first statement is coefficient stabilization in the finite products. The second is [theorem 2.2](#) summed over r . $\square$

Formal versions. The corpus proves the finite factorization as Fabius.thueMorseBlockPolynomial_eq_product (ThueMorseGenerating.lean); ThueMorseBooleanCube.lean now proves it over an *arbitrary* commutative ring as Fabius.prod_one_sub_pow_eq_sum_thueMorseSign, as the specialization $u = -1$ of a two-parameter identity

$$\prod_{j < m} (1 + u z^{2^j}) = \sum_{n < 2^m} u^{\text{wt}_2(n)} z^n \quad (\text{Fabius.prod_one_add_mul_pow}),$$

valid in any commutative semiring with no convergence hypothesis. Its reindexing and digit engine is centralized in ThueMorseBitSupport.lean: the bounded Boolean-cube equivalence Fabius.powersetRangeEquivFin, its sum form Fabius.sum_powerset_two_pow, and the dictionary Fabius.binaryWeight_sum_two_pow, Fabius.sum_two_pow_lt_two_pow. These are exactly the Boolean-cube structures used by the master product.

6.3 Reciprocity and cyclotomic multiplicities

The reflection symmetry becomes a reciprocal-polynomial identity.

Proposition 6.4 (Reciprocity). *For every $m \geq 0$,*

$$z^{2^m-1} P_m(z^{-1}) = (-1)^m P_m(z). \quad (6.6)$$

Thus P_m is reciprocal for even m and anti-reciprocal for odd m .

Proof. Reverse the coefficients and apply (2.8); alternatively, apply $z^{2^j}(1 - z^{-2^j}) = z^{2^j} - 1 = -(1 - z^{2^j})$ to every factor. $\square$

Repeatedly factor $1 - z^{2^j}$ by difference of squares:

$$P_m(z) = (1 - z)^m \prod_{q=0}^{m-2} (1 + z^{2^q})^{m-1-q}. \quad (6.7)$$

If Φ_d is the d -th cyclotomic polynomial, then equivalently

$$P_m(z) = (-1)^m \prod_{q=0}^{m-1} \Phi_{2^q}(z)^{m-q}. \quad (6.8)$$

Here $\Phi_1(z) = z - 1$. In particular, a primitive 2^q -th root of unity is a zero of multiplicity $m - q$ when $0 \leq q < m$. This multiplicity table simultaneously explains Prouhet cancellation at $z = 1$ and the systematic zeros in the dyadic discrete Fourier transform.

Formal versions. `ThueMorseBlockAlgebra.lean` proves the reciprocity (6.6) in two forms: as the coefficient-word reversal `Fabius.reverse_thueMorseBlockPolynomial` ($P_m^{\text{rev}} = (-1)^m P_m$ through `Polynomial.reverse`) and as the evaluation identity `Fabius.sum_thueMorseSign_mul_pow_reflect` over an arbitrary commutative ring, which is (6.6) cleared of the denominator z^{2^m-1} . The degree data feeding it — $\deg P_m = 2^m - 1$, top and leading coefficient $(-1)^m$, $P_m \neq 0$ — are `Fabius.natDegree_thueMorseBlockPolynomial`, `Fabius.coeff_thueMorseBlockPolynomial_top`, `Fabius.leadingCoeff_thueMorseBlockPolynomial`, and `Fabius.thueMorseBlockPolynomial_ne_zero`. The difference-of-squares ladder (6.7) is `Fabius.prod_one_sub_two_pow_eq_ladder`, valid over any commutative ring and built on the telescope $(1 - z) \prod_{i < j} (1 + z^{2^i}) = 1 - z^{2^j}$ (`Fabius.one_sub_mul_prod_one_add_two_pow`); the resulting Prouhet divisibility $(1 - X)^m \mid P_m$ is `Fabius.one_sub_X_pow_dvd_thueMorseBlockPolynomial`. The cyclotomic table (6.8) is `Fabius.thueMorseBlockPolynomial_eq_cyclotomic_prod`, resting on $\Phi_{2^{q+1}} = 1 + z^{2^q}$ (`Fabius.cyclotomic_two_pow_succ`, over any commutative ring).

6.4 The infinite product and Mahler equation

For $|z| < 1$, define

$$E(z) = \sum_{n \geq 0} \varepsilon_n z^n, \quad T(z) = \sum_{n \geq 0} \tau(n) z^n. \quad (6.9)$$

Theorem 6.5 (Infinite lacunary product). *For $|z| < 1$, and coefficientwise as a formal power series,*

$$E(z) = \prod_{j=0}^{\infty} (1 - z^{2^j}). \quad (6.10)$$

The product converges uniformly on compact subsets of the unit disk. The zero–one series is

$$T(z) = \frac{1}{2(1 - z)} - \frac{1}{2} E(z). \quad (6.11)$$

Proof. For $|z| \leq \rho < 1$, the sum $\sum_j |z|^{2^j}$ converges, so the product converges uniformly on the compact disk. Every coefficient stabilizes in the finite products P_m , where [theorem 6.2](#) identifies it as ε_n . The formula for T follows from $\tau(n) = (1 - \varepsilon_n)/2$. $\square$

Peeling off factors gives the Mahler and dissection equations.

Corollary 6.6 (Mahler equations and dissections). *For $|z| < 1$,*

$$E(z) = (1 - z)E(z^2), \quad (6.12)$$

$$E(z) = P_m(z)E(z^{2^m}), \quad (6.13)$$

$$T(z) = (1 - z)T(z^2) + \frac{z}{1 - z^2}. \quad (6.14)$$

Proof. Separate the first factor, or the first m factors, in (6.10). Substitute (6.11) for the final identity; it also follows directly from (2.2). $\square$

Formal versions. `ThueMorseMahler.lean` proves the Mahler equation (6.12) exactly, as an identity of formal power series under `PowerSeries.expand` (`Fabius.thueMorseSeries_mahler`), and the dissection (6.13) at every level with the substitution realized as the m -th iterate of $z \mapsto z^2$ (`Fabius.thueMorseSeries_dissection`, `Fabius.coe_thueMorseBlockPolynomial_series`). The bridge (6.11) is formal in the cleared form $(1 - z)(2T(z) + E(z)) = 1$ (`Fabius.one_sub_X_mul_bridge`: the combination $2T + E$ is the geometric series). The coefficientwise content of theorem 6.5 has long been formal as the stabilization lemma `Fabius.coeff_finite_thueMorse_product` (`ThueMorseGenerating.lean`), and its analytic identity is now formal on the whole complex unit disc. For every $z \in \mathbb{C}$ with $\|z\| < 1$, the signed series converges absolutely (`Fabius.summable_thueMorseSign_mul_pow_complex`), the lacunary product is multipliable (`Fabius.multipliable_one_sub_pow_two_pow_complex`), and

$$\sum_{n \geq 0} \varepsilon_n z^n = \prod_{j \geq 0} (1 - z^{2^j}) \quad (\|z\| < 1), \quad (6.15)$$

by `Fabius.tsum_thueMorseSign_mul_pow_complex` (`ThueMorseInfiniteProduct.lean`). The real APIs `Fabius.summable_thueMorseSign_mul_pow`, `Fabius.multipliable_one_sub_pow_two_pow`, and `Fabius.tsum_thueMorseSign_mul_pow` are retained as specializations and compatibility forms on $-1 < x < 1$, with negative x and $x = 0$ included, together with the unordered and partial-product restatements `Fabius.hasProd_one_sub_pow_two_pow` and `Fabius.tendsto_prod_one_sub_pow_two_pow`. Both sides restrict at scale 2^m to the finite identity (6.3) – valid over any commutative ring – and uniqueness of limits closes the argument. The exponential ray $\sum \varepsilon_n e^{-nt} = \mathcal{E}(t)$ for $t > 0$, which the Dirichlet–Mellin layer consumes, is now the one-line specialization $x = e^{-t}$ (`Fabius.tsum_thueMorseSign_exp_eq_lacunaryExpProduct`). The holomorphic sum `Fabius.thueMorseDiscSeries` is identified with the same complex product by `Fabius.thueMorseDiscSeries_eq_tprod`; its Mahler equation and finite-level factorization $F(z) = P_k(z)F(z^{2^k})$ are `Fabius.thueMorseDiscSeries_mahler` and `Fabius.thueMorseDiscSeries_iterate` (`ThueMorseDiscSeries.lean`), the forms used by the boundary theory. The bit Mahler equation (6.14) and the lacunary quotient are now formal in `ThueMorseLacunaryQuotient.lean`: the quotient $z/(1 - z^2)$ is realized as zG_1 , where the dyadic geometric series $G_j = \sum_k z^{k2^j}$ is proved inverse to $1 - z^{2^j}$ (`Fabius.one_sub_X_pow_mul_dyadicGeometric`) – so that (6.14) becomes the denominator-free identity $T = (1 - z)T(z^2) + zG_1$ (`Fabius.thueMorseBitZSeries_mahler`), and theorem 6.7 holds at every finite level as $T = \sum_{j < M} P_j z^{2^j} G_{j+1} + P_M T(z^{2^M})$ (`Fabius.thueMorseBitZSeries_lacunary_quotient`), by iterating the Mahler step. In the formal statement the last term uses the M -fold `PowerSeries.expand` of T , and `Fabius.one_sub_X_pow_mul_dyadicGeometric` is the bridge that justifies reading G_{j+1} as $(1 - X^{2^{j+1}})^{-1}$. This finite identity is pure formal-power-series algebra; only the analytic passage $M \rightarrow \infty$, including convergence of the remainder to zero for $|z| < 1$, remains informal.

Iterating the inhomogeneous bit equation yields a useful lacunary sum.

Theorem 6.7 (Lacunary quotient expansion). *Let $P_0(z) = 1$ and $P_j(z) = \prod_{h=0}^{j-1} (1 - z^{2^h})$ for $j \geq 1$. Then, for $|z| < 1$,*

$$T(z) = \sum_{j=0}^{\infty} P_j(z) \frac{z^{2^j}}{1 - z^{2^{j+1}}}. \quad (6.16)$$

Proof. After M iterations of (6.14),

$$T(z) = \sum_{j=0}^{M-1} P_j(z) \frac{z^{2^j}}{1 - z^{2^{j+1}}} + P_M(z) T(z^{2^M}).$$

On compact subdisks P_M is bounded and $T(z^{2^M}) \rightarrow T(0) = 0$. $\square$

6.5 A weighted digit enumerator

Both E and the geometric series are specializations of one umbrella product that keeps the whole digit sum, not only its parity.

Theorem 6.8 (Weighted enumerator). *For $|z| < 1$ and any scalar u , and coefficientwise as formal series,*

$$F_u(z) := \sum_{n \geq 0} u^{\text{wt}_2(n)} z^n = \prod_{j \geq 0} (1 + u z^{2^j}), \quad (6.17)$$

with the Mahler equation

$$F_u(z) = (1 + uz) F_u(z^2). \quad (6.18)$$

The specializations $u = 1$ and $u = -1$ are

$$F_1(z) = \frac{1}{1 - z}, \quad F_{-1}(z) = E(z). \quad (6.19)$$

Proof. Expanding the product chooses a finite set of bit positions; by uniqueness of binary expansions each exponent n arises exactly once, with coefficient $u^{\text{wt}_2(n)}$. Peeling the first factor gives (6.18). At $u = 1$ every n receives coefficient one, which restates the uniqueness of binary expansions; at $u = -1$ the coefficient is ε_n . $\square$

Differentiating (6.17) with respect to u at $u = 1$ gives the ordinary generating function of the digit sum itself:

$$\sum_{n \geq 0} \text{wt}_2(n) z^n = \frac{1}{1 - z} \sum_{j \geq 0} \frac{z^{2^j}}{1 + z^{2^j}}. \quad (6.20)$$

Higher derivatives at $u = 1$ generate the falling-factorial moments $\text{wt}_2(n)(\text{wt}_2(n) - 1) \cdots (\text{wt}_2(n) - r + 1)$ in the same way. The lacunary sum in (6.20) is the characteristic-zero twin of the Artin–Schreier solution (7.16) below: the same series $\sum_j z^{2^j} / (1 + z^{2^j})$ computes the digit sum with integer coefficients and the digit parity in characteristic two. The umbrella product itself is formalized in full generality as `Fabius.prod_one_add_mul_pow (ThueMorseBooleanCube.lean)`; the specializations (6.19) are the $z = 1$ and $u = -1$ readings recorded there as `Fabius.sum_pow_binaryWeight_eq_one_add_pow` and `Fabius.prod_one_sub_pow_eq_sum_thueMorseSign`.

6.6 Thue–Morse constants in every base

For a real $b > 1$, put

$$\Theta_b = \sum_{n \geq 0} \frac{\tau(n)}{b^{n+1}}. \quad (6.21)$$

For integer $b \geq 2$, these are the base- b numbers whose fractional digits are the Thue–Morse bits.

Corollary 6.9 (Base- b product). *For every real $b > 1$,*

$$\Theta_b = \frac{1}{2(b-1)} - \frac{1}{2b} \prod_{j \geq 0} (1 - b^{-2^j}). \quad (6.22)$$

The signed companion is

$$\sum_{n \geq 0} \frac{\varepsilon_n}{b^{n+1}} = \frac{1}{b} \prod_{j \geq 0} (1 - b^{-2^j}). \quad (6.23)$$

Proof. Set $z = 1/b$ in (6.11) and (6.10), then multiply by $1/b$. $\square$

Formal versions. Both displays are in `BaseBProduct.lean`, for every real $b > 1$: the signed companion (6.23) is `Fabius.tsum_thueMorseSign_div_pow_succ` and the boxed (6.22) is `Fabius.tsum_thueMorseBit_div_pow_succ`, each stated with the genuine infinite sum $\sum'$ and infinite product $\prod'$. Both are the real infinite product `Fabius.tsum_thueMorseSign_mul_pow` of `ThueMorseInfiniteProduct.lean` at $x = 1/b$, the second after writing $\tau(n) = (1 - \varepsilon_n)/2$ and summing the geometric series.

For $b = 2$, $\Theta_2 = 0.0110100110010110\dots_2$ is the classical Prouhet–Thue–Morse constant. Mahler’s theory implies its transcendence; that deep arithmetic result is separate from the elementary product identity proved here [8, 12].

7 Euler transforms, coefficient recurrences, and algebraic series

7.1 The logarithmic ruler transform

Take the formal logarithm of (6.10):

$$\log E(z) = \sum_{j \geq 0} \log(1 - z^{2^j}) = - \sum_{j \geq 0} \sum_{r \geq 1} \frac{z^{r2^j}}{r}. \quad (7.1)$$

For a fixed exponent k , the admissible j ’s satisfy $0 \leq j \leq \nu_2(k)$, and the corresponding denominator is $k/2^j$. Define

$$a_k = \sum_{j=0}^{\nu_2(k)} 2^j = 2^{\nu_2(k)+1} - 1. \quad (7.2)$$

Theorem 7.1 (Euler transform of the ruler function). *As a formal series, and analytically for $|z| < 1$,*

$$\log E(z) = - \sum_{k \geq 1} \frac{a_k}{k} z^k, \quad a_k = 2^{\nu_2(k)+1} - 1. \quad (7.3)$$

This representation contains no explicit binary digit sum. The coefficient sequence (a_k) is a simple transform of the ruler sequence.

Formal version. [Theorem 7.1](#) is formal in its denominator-free form. A formal logarithm carries the denominators $1/k$ and so lives over $\mathbb{Q}$; differentiating (7.3) and clearing E gives the equivalent identity

$$z E'(z) = -A(z) E(z), \quad A(z) = \sum_{k \geq 1} a_k z^k,$$

which holds over $\mathbb{Z}$, and that is `Fabius.X_mul_derivative_thueMorseSeries` in `RulerLogDerivative.lean`, an identity of formal power series with integer coefficients, with A as `Fabius.rulerSeries` built on the ruler coefficient `Fabius.rulerCoeff` of `ThueMorseHessenberg.lean`. Its coefficient of z^n is exactly the ruler-convolution recurrence of [theorem 7.2](#), from which it is derived; no convergence enters. The two readings the theorem also offers – the $\mathbb{Q}$ -logarithm itself, and the analytic identity for $|z| < 1$ – are not formalized: the first is a coefficientwise integration of the display above, the second its evaluation.

7.2 A nonlocal convolution recurrence

Theorem 7.2 (Ruler-convolution recurrence). *With $\varepsilon_0 = 1$, for every $n \geq 0$,*

$$n\varepsilon_n = - \sum_{k=1}^n a_k \varepsilon_{n-k}. \quad (7.4)$$

Proof. Differentiate (7.3), multiply by $E(z)$, and compare the coefficient of z^{n-1} :

$$E'(z) = -E(z) \sum_{k \geq 1} a_k z^{k-1}.$$

□

The weighted convolution on the right always collapses to $\pm n$. As an evaluator it is inferior to popcount, but it is the natural recurrence attached to the Euler product.

Remark 7.3 (Combinatorial proof of the recurrence). [Theorem 7.2](#) also has a purely finite proof which explains why the convolution always collapses to $\pm n$. Write $a_k = \sum_{j \geq 0, 2^j | k} 2^j$ as the sum of the binary divisors of k and interchange the two sums:

$$\sum_{k=1}^n a_k \varepsilon_{n-k} = \sum_{j \geq 0} 2^j \sum_{\substack{1 \leq k \leq n \\ 2^j | k}} \varepsilon_{n-k} = \sum_{j \geq 0} 2^j \sum_{\substack{0 \leq m < n \\ m \equiv n \pmod{2^j}}} \varepsilon_m.$$

By block concatenation $\varepsilon_{2^j t + r} = \varepsilon_t \varepsilon_r$ ($r < 2^j$), the residue-class sum equals $\varepsilon_{n \bmod 2^j} \sum_{t < \lfloor n/2^j \rfloor} \varepsilon_t$, and the prefix sum collapses by pair cancellation to $\mathbf{1}_{\lfloor n/2^j \rfloor \text{ odd}} \varepsilon_{\lfloor n/2^j + 1 \rfloor}$. When the j -th binary digit of n is 1, the odd-successor law $\varepsilon_{2u+1} = -\varepsilon_u$ turns the product back into $-\varepsilon_n$; when it is 0, the class sum vanishes. Hence $\sum_k a_k \varepsilon_{n-k} = -\varepsilon_n \sum_{j: b_j(n)=1} 2^j = -n\varepsilon_n$: the surviving binary divisors reassemble the binary expansion of n .

Formal versions. `ThueMorseEulerTransform.lean` proves [theorem 7.2](#) as `Fabius.ruler_convolution`, following [remark 7.3](#) – no logarithm, derivative, or convergence enters, and no side condition on n is imposed: the Lean statement is hypothesis-free, the case $n = 0$ being the empty sum. The reusable pieces are the prefix collapse `Fabius.sum_thueMorseSign_range`, the residue-class collapse `Fabius.sum_thueMorseSign_progression`, the divisor-sum identification `Fabius.sum_ite_two_pow_dvd`, the binary expansion `Fabius.sum_ite_odd_div_two_pow`, and the fully general reindexing `Fabius.sum_filter_dvd_sub_eq_sum_progression`, valid for any function into any additive commutative monoid and any modulus. The Euler transform (7.3) itself remains analytic; the partition formula below is now fully formal.

7.3 Partition and Bell-polynomial formulae

Exponentiating (7.3) gives a finite partition sum for every coefficient.

Theorem 7.4 (Partition formula). *For $n \geq 0$,*

$$\varepsilon_n = \sum_{\substack{m_1, \dots, m_n \geq 0 \\ m_1 + 2m_2 + \dots + nm_n = n}} \prod_{j=1}^n \frac{1}{m_j!} \left(-\frac{a_j}{j} \right)^{m_j}. \quad (7.5)$$

Proof. Write $E(z) = \prod_{j \geq 1} \exp(-a_j z^j / j)$, expand each exponential, and collect terms of total degree n . $\square$

Let $\mathcal{B}_n^{\text{exp}}$ denote the complete exponential Bell polynomial, defined by

$$\exp \left(\sum_{j \geq 1} x_j \frac{z^j}{j!} \right) = \sum_{n \geq 0} \mathcal{B}_n^{\text{exp}}(x_1, \dots, x_n) \frac{z^n}{n!}. \quad (7.6)$$

Corollary 7.5 (Bell-polynomial coefficient). *For every $n \geq 0$,*

$$\varepsilon_n = \frac{1}{n!} \mathcal{B}_n^{\text{exp}}(-0!a_1, -1!a_2, \dots, -(n-1)!a_n). \quad (7.7)$$

Proof. Use $x_j = -(j-1)!a_j$ in (7.6). $\square$

Although the right sides are rational-looking sums, they always equal ± 1 . This integrality collapse is itself encoded by the original lacunary product.

Formal versions. [Theorem 7.4](#) is `Fabius.thueMorse_partition_formula` and [corollary 7.5](#) is `Fabius.thueMorse_bell_formula`; their `_field` companions hold over every characteristic-zero field. These Thue–Morse specializations remain in `ThueMorsePartition.lean`.

The reusable finite combinatorics has been extracted to `ExponentialPartition.lean`. For a coefficient sequence $E: \mathbb{N}_0 \rightarrow R$ in any commutative $\mathbb{Q}$ -algebra, `Fabius.partitionExpSum` uses rational reciprocal factorials acting by scalars and satisfies both the marked-part recurrence

$$nP_E(n) = \sum_{j=1}^n j E_j P_E(n-j)$$

(`Fabius.partitionExpSum_recurrence`) and its normalized successor form (`Fabius.partitionExpSum_succ`). No cancellation in R is used, so rings with zero divisors are allowed. `Fabius.partitionExpSum_eq_sum_div` recovers the familiar quotient notation over characteristic-zero fields.

Finally, `ExponentialBell.lean` proves the function equality

$$\text{Fabius.partitionExpSum } E = \text{Fabius.SaddleExpansion.expCoeff } E$$

over the same generality. Thus the recursive all-orders saddle coefficient engine has a finite weighted-partition expansion. The later module `MomentCumulantAlgebra.lean` factorially normalizes that layer into the named transforms `Fabius.completeBellPolynomial` and `Fabius.momentCumulant`, proves the classical Bell recurrence and their normalized inverse laws, and establishes functoriality over commutative $\mathbb{Q}$ -algebras. The downstream module `PolynomialExpectationCumulant.lean` applies these transforms to the integral of an arbitrary finite polynomial, first as a raw-moment sum and then as a complete-Bell sum with the exact zeroth-mass correction; its normalized and probability-measure forms are corollaries. This is a generic finite polynomial interface, not a Thue–Morse-specific law or an infinite analytic expansion.

For $E_j = -a_j/j$, the ruler convolution ([theorem 7.2](#)) says that the sign sequence obeys the same recurrence; the division-by- n uniqueness lemma `Fabius.eq_of_ruler_recurrence` then identifies the two. No power series, exponentials, or convergence enter this specialization.

7.4 A Hessenberg determinant

For $n \geq 0$, let

$$H_n = \begin{pmatrix} a_1 & 1 & 0 & \cdots & 0 \\ a_2 & a_1 & 2 & \ddots & \vdots \\ a_3 & a_2 & a_1 & \ddots & 0 \\ \vdots & \vdots & \ddots & \ddots & n-1 \\ a_n & a_{n-1} & \cdots & a_2 & a_1 \end{pmatrix}. \quad (7.8)$$

For $n = 0$ this is the empty matrix, with $\det H_0 = 1$.

Theorem 7.6 (Determinant formula). *For every $n \geq 0$,*

$$\varepsilon_n = \frac{(-1)^n}{n!} \det H_n. \quad (7.9)$$

In particular,

$$\det H_n = (-1)^n n! \varepsilon_n \in \{n!, -n!\}, \quad |\det H_n| = n!. \quad (7.10)$$

Proof. Put $d_0 = 1$ and $d_n = \det H_n$. Expansion of a lower Hessenberg determinant along its first column gives

$$d_n = \sum_{k=1}^n (-1)^{k-1} \frac{(n-1)!}{(n-k)!} a_k d_{n-k}. \quad (7.11)$$

Set $d_n = (-1)^n n! c_n$. Then (7.11) reduces to

$$n c_n = - \sum_{k=1}^n a_k c_{n-k},$$

with $c_0 = 1$. By uniqueness of the recurrence in [theorem 7.2](#), $c_n = \varepsilon_n$. $\square$

Remark 7.7 (An expansion-free proof). The formalization (`Fabius.det_thueMorseHessenberg`, `ThueMorseHessenberg.lean`) avoids the Hessenberg expansion (7.11) altogether. Write the ruler convolution for $\varepsilon(1), \dots, \varepsilon(n)$ as a linear system $T_n \mathbf{x} = \mathbf{c}$ whose matrix T_n is lower triangular with diagonal $1, 2, \dots, n$, so $\det T_n = n!$ (`Fabius.det_thueMorseTriangular`). Cramer’s rule – in the reusable form $\text{cramer}_A(A\mathbf{x}) = \det A \cdot \mathbf{x}$, proved for any commutative ring (`Matrix.cramer_mulVec`, `ToeplitzHessenberg.lean`) – evaluates $n! \varepsilon(n)$ as the determinant of T_n with its last column replaced by $\mathbf{c}$; one cyclic column rotation ($\text{sign } (-1)^{n-1}$) and one column negation turn that matrix into H_n . Hence $\det H_n = (-1)^n n! \varepsilon(n)$ with no minor ever expanded, for every $n \geq 0$; that $\det H_n = \pm n!$ is `Fabius.det_thueMorseHessenberg_eq_or`, and the sharper $|\det H_n| = n!$ is `Fabius.natAbs_det_thueMorseHessenberg`.

None of this is special to the ruler coefficients. The argument is carried out in `ToeplitzHessenberg.lean` for arbitrary band data (d_m) and (a_k) over an arbitrary commutative ring: the triangular matrix `Fabius.triBand`, of size N and with $\det = \prod_{i < N} d_i$ (`Fabius.det_triBand`); the lower Toeplitz–Hessenberg matrix `Fabius.hessBand`; and the determinant formula $\det = (-1)^{n+1} (\prod_{i \leq n} d_i) x_n$ for a matrix of size $n+1$, read off from *any* solution x of the convolution system $T\mathbf{x} = -(a_1, \dots, a_{n+1})$ (`Fabius.det_hessBand`), together with the scalar form `Fabius.det_hessBand_of_recurrence` that the Thue–Morse instance consumes. No invertibility, no positivity, and no uniqueness of the solution is assumed anywhere.

7.5 The zero–one series over $\mathbb{F}_2$

The signed series loses all information modulo two because $-1 = 1$. The bit series remains nontrivial. Work in $\mathbb{F}_2[[z]]$ and write

$$\Theta(z) = \sum_{n \geq 0} \tau(n) z^n. \quad (7.12)$$

Theorem 7.8 (Artin–Schreier equation). *In $\mathbb{F}_2[[z]]$,*

$$\Theta(z) = (1+z)\Theta(z^2) + \frac{z}{1+z^2}, \quad (7.13)$$

$$\boxed{(1+z)^3\Theta(z)^2 + (1+z)^2\Theta(z) + z = 0.} \quad (7.14)$$

If $Y(z) = (1+z)\Theta(z)$, then

$$\boxed{Y(z)^2 + Y(z) = \frac{z}{1+z}.} \quad (7.15)$$

Proof. Split the series into even and odd coefficients and use $\tau(2n) = \tau(n)$, $\tau(2n+1) = 1 + \tau(n)$ in characteristic two:

$$\Theta(z) = \Theta(z^2) + z \sum_{n \geq 0} (1 + \tau(n)) z^{2n} = (1+z)\Theta(z^2) + \frac{z}{1+z^2}.$$

Frobenius gives $\Theta(z)^2 = \Theta(z^2)^2$, while $1 + z^2 = (1+z)^2$. Clear the denominator and substitute $Y = (1+z)\Theta$. $\square$

Corollary 7.9 (Explicit Artin–Schreier sum). *The unique zero-constant-term solution is*

$$Y(z) = \sum_{j=0}^{\infty} \left(\frac{z}{1+z} \right)^{2^j} = \sum_{j=0}^{\infty} \frac{z^{2^j}}{1+z^{2^j}}, \quad (7.16)$$

$$\boxed{\Theta(z) = \frac{1}{1+z} \sum_{j=0}^{\infty} \frac{z^{2^j}}{1+z^{2^j}}.} \quad (7.17)$$

Proof. Let $A = z/(1+z)$. In the z -adic topology, $Y = \sum_{j \geq 0} A^{2^j}$ is well defined and $Y^2 + Y = A$ by index shifting. Frobenius gives $A^{2^j} = z^{2^j}/(1+z^{2^j})$. A zero-constant solution of the homogeneous equation $Z^2 + Z = 0$ must vanish, proving uniqueness. $\square$

Christol’s theorem says that a finite-field sequence is automatic exactly when its generating series is algebraic. Here the degree-two certificate is derived directly; it is an algebraic shadow of the two-state automaton [13].

Formal versions. `ThueMorseArtinSchreier.lean` proves the functional equation (7.13) in the cleared form $\Theta = (1+z)\Theta^2 + zG^2$ with $G = \sum z^n$ (`Fabius.thueMorseBitSeries_functional`, `Fabius.one_add_X_mul_geometricSeries`), the boxed quadratic (7.14) as `Fabius.thueMorseBitSeries_quadratic`, and the normalized form (7.15) cleared of its denominator, $(1+z)(Y^2 + Y) = z$, as `Fabius.artinSchreier_one_add_X_mul`. The Frobenius step $\Theta(z)^2 = \Theta(z^2)^2$ is isolated as a coefficientwise lemma valid over *any* commutative ring of characteristic two (`Fabius.coeff_sq_charTwo`: the off-diagonal convolution terms cancel in pairs), specializing over $\mathbb{F}_2$ to `Fabius.coeff_sq_zmodTwo`. Of the explicit solution (7.16) the formalization keeps the finite content: the partial sums $S_J = \sum_{j < J} a^{2^j}$ satisfy $S_J^2 + S_J = a^{2^J} + a$ in any commutative ring of characteristic two (`Fabius.sum_pow_two_pow_sq_add`), so they solve the Artin–Schreier equation up to the single tail term, and no z -adic limit is needed for any identity in this subsection.

The quadratic’s *entire* solution set is now formal as well, and it is a two-element set: in $\mathbb{F}_2[[z]]$,

$$(1+z)^3Y^2 + (1+z)^2Y + z = 0 \iff Y = \Theta \text{ or } Y = \Theta + G \quad (7.18)$$

(`Fabius.artinSchreier_solution_iff`), with no hypothesis on Y whatever — in particular none on its constant term. The difference $D = Y - \Theta$ of two solutions satisfies the factored identity $(1+z)^2D((1+z)D+1) = 0$; since $\mathbb{F}_2[[z]]$ is a domain and $1+z \neq 0$, either $D = 0$ or $(1+z)D = 1$, and the second alternative is not a contradiction but the second root: it says $D = G$, because $(1+z)G = 1$. The Artin–Schreier symmetry $Y \mapsto Y + 1$ of (7.15) is exactly the translation $\Theta \mapsto \Theta + G$ after the substitution $Y = (1+z)\Theta$, and $\Theta + G$ is verified to be a root independently (`Fabius.thueMorseBitSeries_add_geometricSeries_quadratic`).

7.6 An integral algebraic lift

The fourth draft records a rather different integer sequence whose parity is Thue–Morse. Define $c_0 = 0$, $c_1 = 1$, and for $n \geq 2$,

$$nc_n = (5 - 2n)c_{n-1} + 3(n-1)c_{n-2} + 1. \quad (7.19)$$

Its first terms are

$$0, 1, 1, 2, 1, 4, -2, 13, -23, 68, -164, 439, -1146, 3067, \dots$$

Theorem 7.10 (Integer algebraic lift). *Let $C(z) = \sum_{n \geq 0} c_n z^n$. Then*

$$C(z) = \frac{1}{2(1-z)} \left(\sqrt{\frac{1+3z}{1-z}} - 1 \right), \quad (7.20)$$

$$(1-z)^3 C(z)^2 + (1-z)^2 C(z) = z. \quad (7.21)$$

All c_n are integers and

$$\boxed{c_n \equiv \tau(n) \pmod{2}, \quad \varepsilon_n = (-1)^{c_n}.} \quad (7.22)$$

Proof. Multiplying (7.19) by z^{n-1} , summing, and including the initial terms gives a first-order differential equation whose zero-constant solution is (7.20); squaring yields (7.21).

Conversely, the coefficient of z^n in $(1-z)^2 C(z)$ contains c_n with coefficient one, while the quadratic term uses only earlier coefficients because $C(0) = 0$. Thus (7.21) recursively determines integers c_n .

Modulo two, (7.21) becomes precisely (7.14). That equation has a unique zero-constant solution in $\mathbb{F}_2[[z]]$, so $C(z) \equiv \Theta(z) \pmod{2}$. $\square$

Remark 7.11. The lift grows and changes sign; it is not a disguised digit sum. Its parity is Thue–Morse because an integral algebraic equation reduces to the automatic sequence’s Artin–Schreier equation. This suggests a reusable discovery method: search for integral algebraic lifts of finite-field automatic series.

Remark 7.12 (Closed form and a purely arithmetic parity proof). Expanding (7.20) with $u = z/(1-z)$ against the Catalan generating function gives the explicit coefficients

$$c_n = \sum_{k=1}^n (-1)^{k-1} C_{k-1} \binom{n}{k}, \quad (7.23)$$

where C_m is the m -th Catalan number — manifestly integers, with no recursion needed. This representation yields a parity proof that bypasses power series entirely: Catalan numbers are odd exactly at the Mersenne indices $m = 2^j - 1$ (fold the convolution $C_{m+1} = \sum C_i C_{m-i}$ at its middle and recurse), so modulo two the sum (7.23) retains only $k = 2^j$; by Lucas’s theorem $\binom{n}{2^j}$ is odd exactly when the j -th bit of n is set; and the bits sum to $\text{wt}_2(n)$. Hence $c_n \equiv \text{wt}_2(n) \equiv \tau(n) \pmod{2}$.

Formal versions. The uniqueness step is `Fabius.artinSchreier_solution_unique` (`ThueMorseArtinSchreier.lean`), and it is now a one-line consequence of the complete solution set (7.18): the quadratic has exactly the two roots Θ and $\Theta + G$, and of these only Θ has constant term zero, because G has constant term one. The statement itself is unchanged — a zero-constant-term solution of the quadratic equals Θ — only its proof now factors through the stronger classification. The parity statement (7.22) is formal in conditional-universal form: *any* integer sequence with $c_0 = 0$ whose generating series satisfies (7.21) has Thue–Morse parity (`Fabius.integerLift_parity`, `Fabius.integerLift_modEq`). The sequence itself is now *constructed* in `ThueMorseIntegerLift.lean` via the closed form (7.23) (`Fabius.integerLift`), and the boxed parity claims are proved unconditionally along the route of [remark 7.12](#): Catalan parity is `Fabius.catalan_odd_iff` (via `Fabius.catalan_mod_two`, reusable facts about C_m independent of this atlas), and $c_n \equiv \tau(n) \pmod{2}$, $\varepsilon_n = (-1)^{c_n}$ are `Fabius.integerLift_cast_zmod`, `Fabius.integerLift_emod_two`, and `Fabius.thueMorseSign_eq_neg_one_pow_integerLift`. The algebraic equation (7.21) is now proved for the constructed series (`Fabius.integerLiftSeries_algebraic`, `ThueMorseIntegerLiftEquation.lean`) by pure finite combinatorics: the second difference of the lift equals the substituted alternating Catalan series (`Fabius.integerLift_delta_bridge`, double Pascal), the interior product sum collapses through a reusable *diagonal Vandermonde* identity (`Fabius.sum_range_choose_mul_choose`), the substituted series satisfies the quadratic convolution identity $\sum_{p+q=m} h_p h_q = h_m - h_{m+1} - [m=0]$ (`Fabius.catalanSeriesDelta_conv`) — the coefficient form of $H = 1 - uH^2$ — and cancelling the non-zero-divisor $1 - z$ in the domain $\mathbb{Z}[[z]]$ yields the boxed equation. Finally, the defining recurrence (7.19) is `Fabius.integerLift_recurrence` (`ThueMorseIntegerLiftRecurrence.lean`): differentiating the algebraic equation and linearizing with one certificate produces the first-order equation $(1 + 3z)(1 - z)^2 C' = 1 + 3(1 - z^2)C$ (`Fabius.integerLiftSeries_ode`), whose coefficients give a four-term relation that telescopes onto the recurrence by induction. Every clause of [theorem 7.10](#) is now formal and unconditional.

8 Natural boundary, Hankel determinants, and Diophantine consequences

The algebraic picture of the previous section is completed by three analytic-arithmetic facts: the series E cannot be continued across any boundary point, its Hankel determinants never vanish, and its values at algebraic points are transcendental with the smallest possible irrationality exponent. The first two facts are proved here in full — and both are formalized; the Diophantine statements are deep literature theorems, stated with references and not reproved.

8.1 The unit circle is a natural boundary

A *natural boundary* of a power series is a curve through which no analytic continuation across any arc is possible.

Theorem 8.1 (Natural boundary). *Every point of the unit circle is a singular point of E ; the circle is a natural boundary. Consequently E and T are transcendental over the field $\mathbb{C}(z)$ of rational functions.*

Proof. Let ζ be a root of unity of two-power order, say $\zeta^{2^r} = 1$. For $0 < \rho < 1$, split the product (6.10) at index r : the head has r factors, each of modulus at most 2, while every tail factor $1 - (\rho\zeta)^{2^j} = 1 - \rho^{2^j}$ with $j \geq r$ is a real number in $(0, 1)$. Hence

$$|E(\rho\zeta)| \leq 2^r (1 - \rho^{2^r}) \xrightarrow[\rho \uparrow 1]{} 0.$$

Such roots of unity are dense on the circle. If E continued analytically through some open boundary arc, the continued function would vanish on a set with an accumulation point inside its domain, hence

would vanish identically by the identity theorem, contradicting $E(0) = 1$. An algebraic function over $\mathbb{C}(z)$ has only finitely many singular points on the Riemann sphere, so it cannot have a natural boundary; and T differs from $-E/2$ by the rational function in (6.11). $\square$

Formal versions. The natural-boundary core of [Theorem 8.1](#) is fully formal. The series $F(z) = \sum \varepsilon_n z^n$ is realized as a holomorphic function on the open unit disc by locally uniform convergence of its polynomial partial sums (`Fabius.thueMorseDiscSeries`, `Fabius.thueMorseDiscSeries_differentiableOn`, `ThueMorseDiscSeries.lean`), satisfying the analytic Mahler equation $F(z) = (1 - z)F(z^2)$ and its iterate $F(z) = \prod_{j < k} (1 - z^{2^j}) F(z^{2^k})$ (`Fabius.thueMorseDiscSeries_mahler`, `Fabius.thueMorseDiscSeries_iterate`), with the real-ray bound $|F(x)| \leq 1 - x$ on $[0, 1)$ (`Fabius.thueMorseDiscSeries_real_le`) — the first factor of the lacunary product already vanishes at the boundary. The radial vanishing at every two-power root of unity is `Fabius.tendsto_thueMorseDiscSeries_radial_zero`, and the theorem itself is `Fabius.thueMorse_natural_boundary`: no function holomorphic on any disc centered at a point of the unit circle agrees with F on the overlap with the unit disc — the hypothetical extension vanishes at the dyadic roots accumulating at the center, hence identically by the identity theorem, forcing $F \equiv 0$ on the disc against $F(0) = 1$. Only the algebraic coda — that an algebraic function over $\mathbb{C}(z)$ has finitely many singular points, whence the transcendence of E and T over $\mathbb{C}(z)$ — remains informal.

Remark 8.2. Two softer routes to the same conclusion deserve mention. The Pólya–Carlson theorem [20] says that a power series with integer coefficients and radius of convergence one is either rational or has the unit circle as a natural boundary, and E is visibly not rational because its coefficient sequence is not eventually periodic ([section 14](#)). For Mahler functions there is also a general rational-or-natural-boundary dichotomy [21]. The radial-zero argument above is preferred here because it is self-contained.

The contrast with [section 7.5](#) is worth displaying:

$$\left. \begin{array}{l} \text{over } \mathbb{F}_2(z) \\ \text{over } \mathbb{C}(z) \end{array} \right| \begin{array}{l} \Theta(z) \text{ is algebraic of degree two,} \\ T(z) \text{ is transcendental with a natural boundary.} \end{array} \quad (8.1)$$

Automaticity is algebraicity in the matching finite characteristic, not in characteristic zero.

8.2 Hankel determinants never vanish

For a sequence $c = (c_n)_{n \geq 0}$, the order- r *Hankel determinant* is the determinant of the $r \times r$ matrix whose (i, j) -entry is c_{i+j} :

$$H_r(c) = \det(c_{i+j})_{0 \leq i, j < r}. \quad (8.2)$$

For the signed sequence the first values are

$$H_1, H_2, \dots = 1, -2, 4, 8, -16, -32, -64, 128, -256, -1536, \dots \quad (8.3)$$

Unlike the Hessenberg matrix of [theorem 7.6](#), which was engineered from the ruler weights precisely so that its determinant would telescope, these matrices sample the raw signs and have no visible triangular structure.

Theorem 8.3 (Allouche–Peyrière–Wen–Wen). *For every $r \geq 1$,*

$$\boxed{H_r(\varepsilon) \neq 0.} \quad (8.4)$$

Proof. We show the exact two-adic valuation $H_r(\varepsilon) = 2^{r-1} \cdot (\text{odd})$, which is visible in (8.3) — for instance $H_{10} = -1536 = -2^9 \cdot 3$. Interleave the rows and the columns of H_{2n} by parity (the same permutation on both sides, so no sign). By $\varepsilon_{2k} = \varepsilon_k$, $\varepsilon_{2k+1} = -\varepsilon_k$ the four blocks

are $\begin{bmatrix} A & -A \\ -A & K \end{bmatrix}$ with A the order- n Hankel matrix and $K_{ij} = \varepsilon_{i+j+1}$; adding the top block row to the bottom one triangularizes, and the Schur block has entries $\varepsilon_{i+j+1} - \varepsilon_{i+j} = 2m_{i+j}$, where $m_k = \tau(k) - \tau(k+1) \in \{-1, 0, 1\}$ is the half-increment. Hence

$$H_{2n}(\varepsilon) = 2^n H_n(\varepsilon) H_n(m), \quad H_{2n+1}(\varepsilon) = 2^n H_{n+1}(\varepsilon) H_n(m),$$

the odd case by the same elimination against the truncated top block. The determinants $H_n(m)$ are *odd*: modulo two, $|m|$ is the period-doubling sequence p with $p(2k) = 1$, $p(2k+1) = 1 + p(k)$, and the two mod-two determinant families – Hankel determinants $D(n)$ of $p(k)$ and $Y(n)$ of $p(\lfloor k/2 \rfloor)$ – close under the same parity interleaving:

$$D(2n) = D(n), \quad D(2n+1) = Y(n), \quad Y(2n) = D(n)Y(n), \quad Y(2n+1) = D(n+1)Y(n),$$

using $p(k) + p(k+1) = p(\lfloor k/2 \rfloor)$ and the vanishing of all permutation signs in characteristic two. Every value is 1 by strong induction. Feeding the oddness back into the two factorizations gives $\nu_2(H_{2n}) = n + \nu_2(H_n)$ and $\nu_2(H_{2n+1}) = n + \nu_2(H_{n+1})$, whence $\nu_2(H_r) = r - 1$ exactly and $H_r \neq 0$. $\square$

The original proof runs a closed finite system of determinant recurrences [22]; later proofs include Han’s Hankel continued-fraction calculus [23] and a combinatorial argument of Bugeaud and Han [24]. The two-adic argument above is in the spirit of all three, compressed to a pair of block factorizations and one mod-two induction. The statement is a strong rigidity property: no finite window of the sequence ever satisfies an accidental low-order linear relation of Hankel type.

Nonvanishing has a standard approximation-theoretic reading. A *Padé approximant of type* $[r-1/r]$ to E is a rational function $A(z)/B(z)$ with $\deg A \leq r-1$ and $\deg B \leq r$ whose Taylor expansion agrees with E to the maximal possible order. When $H_r(\varepsilon) \neq 0$, this approximant exists and the agreement is exact through order z^{2r-1} with a nonzero discrepancy coefficient at z^{2r} ; the whole diagonal Padé table of E is therefore normal, with no degenerate blocks.

Formal versions. **Theorem 8.3** is fully formal, with the proof above. The mod-two engine is `PeriodDoublingHankel.lean`: the period-doubling bit and its recursions (`Fabius.pdBit`), the four determinant halvings, and the strong induction `Fabius.pdHankel_det_eq_one` – every Hankel determinant of the period-doubling sequence over $\mathbb{F}_2$ equals 1. The reusable block machinery (`Fabius.det_fromBlocks_of_eq_mul`, `Fabius.det_fromBlocks_symm`, the parity interleavings) is proved over an arbitrary commutative ring. The integer level is `ThueMorseHankel.lean`: the doubling factorizations (`Fabius.hankelSign_det_two_mul`, `Fabius.hankelSign_det_two_mul_add_one`), the mod-two transfer (`Fabius.mseqHankel_det_odd`), the exact valuation $H_r = 2^{r-1} \cdot \text{odd}$ (`Fabius.hankelSign_det_eq_two_pow_mul_odd`), and the boxed nonvanishing (8.4) (`Fabius.thueMorseHankel_ne_zero`).

8.3 Transcendental values and the sharp irrationality exponent

Beyond transcendence of the function, Mahler’s method gives transcendence of its values. The Mahler equation (6.12) relates z to z^2 , and for an algebraic α with $0 < |\alpha| < 1$ the forward orbit $\alpha, \alpha^2, \alpha^4, \dots$ stays in the unit disk and never meets the zero of the rational cofactor $1 - z$.

Theorem 8.4 (Mahler; literature theorem). *If α is algebraic and $0 < |\alpha| < 1$, then $E(\alpha)$ is transcendental; consequently so is $T(\alpha)$ [8].*

In particular, every base- b Thue–Morse constant Θ_b of (6.21) is transcendental for integer $b \geq 2$. How well such a constant can be approximated by rationals is measured by its *irrationality exponent* $\mu(\xi)$: the supremum of all real μ such that $|\xi - p/q| < q^{-\mu}$ holds for infinitely many fractions p/q . Every irrational number has $\mu \geq 2$, by Dirichlet’s pigeonhole theorem, and Lebesgue-almost every real number has $\mu = 2$.

Theorem 8.5 (Bugeaud; literature theorem). *For every integer $b \geq 2$,*

$$\boxed{\mu(\Theta_b) = 2.} \tag{8.5}$$

Thus the Thue–Morse–Mahler numbers are transcendental yet approximate as badly as possible for irrational numbers. Bugeaud’s proof [25] runs the Padé approximants of the previous subsection through the Mahler equation: iterating $z \mapsto z^2$ transports one good rational approximation at $z = 1/b$ to a whole geometric family, and the nonvanishing Hankel determinants guarantee that the approximation orders have no long gaps. Earlier partial bounds and the general determinant framework are surveyed in [26, 27]. This is the cleanest place where the determinant rigidity of [theorem 8.3](#) pays an explicitly Diophantine dividend.

Formal versions. Neither theorem of this subsection has one, and neither is planned: both are imported literature results about the *analytic value* of E at an algebraic point, and transcendence and irrationality exponents are outside what the corpus develops. Two algebraic ingredients used across those proofs are formal, but neither proves an analytic evaluation, transcendence statement, or irrationality exponent. The Mahler equation (6.12) is formal-power-series algebra: `Fabius.thueMorseSeries_mahler` in `ThueMorseMahler.lean`, in the form $E = (1 - X) \cdot \text{expand}_2 E$, and the determinant rigidity of [theorem 8.3](#) is `Fabius.thueMorseHankel_ne_zero` in `ThueMorseHankel.lean`, a finite integer-matrix theorem proved through the sharper `Fabius.hankelSign_det_eq_two_pow_mul_odd` — the determinant is a power of two times an odd number, so its nonvanishing is a parity fact rather than an estimate.

9 Finite differences, Prouhet cancellation, and all dyadic moments

9.1 The mixed-difference master identity

For a function f , let

$$(\Delta_h f)(x) = f(x + h) - f(x), \quad (T_h f)(x) = f(x + h). \tag{9.1}$$

Thus $I - T_h = -\Delta_h$.

Theorem 9.1 (Dyadic finite-difference factorization). *For every $m \geq 0$, every scalar h , and every function for which the values are defined,*

$$\boxed{\sum_{n=0}^{2^m-1} \varepsilon_n f(x + nh) = \prod_{j=0}^{m-1} (I - T_{2^j h}) f(x)} \tag{9.2}$$

$$\boxed{= (-1)^m (\Delta_h \Delta_{2h} \cdots \Delta_{2^{m-1}h} f)(x).} \tag{9.3}$$

Proof. Expand the operator product. A subset $S \subseteq \{0, \dots, m-1\}$ contributes the shift by $h \sum_{j \in S} 2^j$ and sign $(-1)^{|S|}$. Binary expansion identifies the subset sum with a unique $n < 2^m$, and the sign is ε_n . $\square$

This theorem is stronger than any one moment identity: it is an inclusion–exclusion formula valid for arbitrary test functions.

Formal versions. `ThueMorseMixedDifference.lean` proves [theorem 9.1](#) in full generality: the argument space is any additive commutative monoid, the value space any additive commutative group, and f an arbitrary function between them. The operator $\prod_{j < m} (I - T_{2^j h})$ is `Fabius.dyadicMixedDifference`, the boxed identity (9.2) is `Fabius.sum_thueMorseSign_smul_eq_mixed`

Difference, and the expanded inclusion–exclusion form is `Fabius.sum_thueMorseSign_smul_eq_sum_powerset_smul`. The polynomial and exponential specializations recover the Prouhet and moment layers formalized earlier.

`ThueMorseSymmetricDifference.lean` now supplies the centered arbitrary-step companion. Its `Fabius.symmetricMixedDifference` uses half-steps a_i : the corresponding single factor is $f(a_i +_v x) - f((-a_i) +_v x)$, where $+_v$ is the additive action, and the full operator is the ordinary mixed difference with steps $-2a_i$ based at the upper cube corner. The general layer assumes only additive commutative groups and an additive action, permits repeated half-step values, and expands over the exact symmetric Boolean cube.

Over a commutative ring, a polynomial p of degree at most $|s|$ is reduced to

$$[X^{|s|}]p \, |s|! \prod_{i \in s} (2a_i).$$

Degree below $|s|$ is annihilated, with the degree-valued theorem retaining the zero-polynomial/empty-support boundary, and the first surviving monomial has the displayed sign-free top factor. For dyadic half-steps the group-valued block is a signed Thue–Morse sum. If the argument is a characteristic-zero field, the increasing-grid specialization is exact for every m , including $m = 0$:

$$D_m(2^{-m}; f, x) = \sum_{n < 2^m} (-1)^m \varepsilon_n \cdot f(x - (1 - 2^{-m}) + 2n2^{-m}).$$

At positive order $N = m + 1$, the full grid point is

$$x - (1 - 2^{-(m+1)}) + \frac{n}{2^m};$$

its varying term is $n/2^m$, and its mesh spacing is $1/2^m$.

The module’s exhaustive public surface is two definitions and eleven theorems:

No.	Declaration	Formal content
1	<code>Fabius.symmetricMixedDifference</code>	Centered arbitrary-half-step definition.
2	<code>Fabius.symmetricMixedDifference_empty</code>	Empty operator.
3	<code>Fabius.symmetricMixedDifference_singleton</code>	One centered factor.
4	<code>Fabius.symmetricMixedDifference_insert</code>	Insert recurrence; half-step values may repeat.
5	<code>Fabius.symmetricMixedDifference_eq_sum_powerset_smul</code>	Symmetric Boolean-cube expansion.
6	<code>Fabius.symmetricMixedDifference_polynomial_eq_coeff_card</code>	Sharp coefficient extraction over a commutative ring.
7	<code>Fabius.symmetricMixedDifference_polynomial_of_degree_lt</code>	Cancellation below the support size.
8	<code>Fabius.symmetricMixedDifference_pow_card</code>	First surviving monomial value.
9	<code>Fabius.symmetricDyadicMixedDifference</code>	Dyadic centered definition.
10	<code>Fabius.symmetricDyadicMixedDifference_zero</code>	Zeroth dyadic operator.
11	<code>Fabius.symmetricDyadicMixedDifference_eq_sum_thueMorseSign_smul</code>	Centered dyadic Thue–Morse block.
12	<code>Fabius.symmetricDyadicMixedDifference_increasing_two_pow_eq_sum_thueMorseSign_smul</code>	Increasing affine grid, including order zero.

No.	Declaration	Formal content
13	<code>Fabius.symmetricDyadicMixedDifference_in_v_two_pow_succ_eq_sum_thueMorseSign_smul</code>	Positive-order grid term $n/2^m$ (mesh $1/2^m$).

This algebraic module does not prove the next multiple-integral corollary. More precisely, it closes the Boolean-cube, arbitrary-step polynomial extraction/cancellation/top-value, and exact affine-grid clauses of the continuous-chaos report’s Thue–Morse corner theorem, but not that theorem’s repeated-integral identity under its C^N hypotheses. The report theorem therefore remains near-complete rather than fully formalized; a final analytic bridge and report-shaped wrapper remain. Its following Walsh conditional-expectation corollary also still lacks the sign–magnitude probability construction, conditioning, and 2^{-N} normalization.

9.2 Exact multiple-integral remainder

Corollary 9.2 (Integral cubature formula). *Suppose $h \geq 0$ and $f \in C^m$ on an interval containing all relevant points. Then*

$$\sum_{n < 2^m} \varepsilon_n f(x + nh) = (-1)^m \int_0^h \int_0^{2h} \cdots \int_0^{2^{m-1}h} f^{(m)}(x + t_0 + \cdots + t_{m-1}) dt_{m-1} \cdots dt_0 \quad (9.4)$$

$$= (-1)^m 2^{\binom{m}{2}} h^m \int_{[0,1]^m} f^{(m)}\left(x + h \sum_{j=0}^{m-1} 2^j u_j\right) d\mathbf{u}. \quad (9.5)$$

Consequently

$$\left| \sum_{n < 2^m} \varepsilon_n f(x + nh) \right| \leq 2^{\binom{m}{2}} h^m \sup_{x \leq y \leq x + (2^m - 1)h} |f^{(m)}(y)|. \quad (9.6)$$

Proof. Apply $(I - T_a)f(x) = -\int_0^a f'(x + t) dt$ successively. Rescale $t_j = 2^j h u_j$ for the unit-cube form. The box volume is $h^m 2^{0+1+\cdots+(m-1)}$. $\square$

Formal version. The nested-integral display (9.4) is `Fabius.thueMorseBlockSum_eq_nestedDyadicIntegral` in `ThueMorseCubature.lean`, for every f of class C^m on the whole line and every real x, h (the sign of h plays no role in the identity itself): the block sum `Fabius.thueMorseBlockSum` equals $(-1)^m$ times `Fabius.nestedDyadicIntegral` of $f^{(m)}$, the nested integral being defined recursively with the outermost variable on $[0, h]$ and the step doubling inward, exactly as displayed. The proof is the induction sketched above, run for all base points at once: splitting the block by its low bit (`Fabius.thueMorseBlockSum_succ`) turns each pair into $-\int_0^h f'(x + 2nh + t) dt$, and the integrand of the outer integral is then the block sum of length 2^m for f' at base point $x + t$ with step $2h$ – a finite sum of continuous functions, so the interchange needs no parametric-integral theory. The bound (9.6) is `Fabius.abs_thueMorseBlockSum_le`, with the sup taken over $[x, x + (2^m - 1)h]$ and the box volume written as $2^{0+1+\cdots+(m-1)} h^m$; the exponent is $\binom{m}{2}$ by `Fabius.sum_range_id_eq_choose_two`. It follows from the volume bound `Fabius.abs_nestedDyadicIntegral_le` on the nested integral. The unit-cube form (9.5) is not formalized.

Corollary 9.3 (Sign for completely monotone kernels). *If $(-1)^m f^{(m)} \geq 0$ on the interval, then*

$$\sum_{n < 2^m} \varepsilon_n f(x + nh) \geq 0. \quad (9.7)$$

For $a, s > 0$,

$$\sum_{n < 2^m} \frac{\varepsilon_n}{(a+n)^s} = (s)_m \int_0^1 \int_0^2 \cdots \int_0^{2^{m-1}} \frac{dt_{m-1} \cdots dt_0}{(a+t_0+\cdots+t_{m-1})^{s+m}} > 0, \quad (9.8)$$

where $(s)_m = s(s+1)\cdots(s+m-1)$. For $m \geq 1$,

$$\sum_{n < 2^m} \varepsilon_n \log(a+n) < 0, \quad 0 < \prod_{n < 2^m} (a+n)^{\varepsilon_n} < 1. \quad (9.9)$$

Proof. Use $f(y) = (a+y)^{-s}$, for which $f^{(m)}(y) = (-1)^m (s)_m (a+y)^{-s-m}$. For the logarithm, $f^{(m)}(y) = (-1)^{m-1} (m-1)! / (a+y)^m$, so the integral is strictly negative. Exponentiation gives the product inequality. $\square$

Formal version. The sign statement (9.7) is `Fabius.thueMorseBlockSum_nonneg` in `ThueMorseCubature.lean`: for $h \geq 0$ and $(-1)^m f^{(m)} \geq 0$ on $[x, x + (2^m - 1)h]$ – the interval that actually contains every argument $x + t_0 + \cdots + t_{m-1}$ – the block sum is nonnegative, by `Fabius.nestedDyadicIntegral_nonneg` applied to $(-1)^m f^{(m)}$ through `Fabius.nestedDyadicIntegral_const_mul`. The *strict* form is formal as well: for $h > 0$ and $(-1)^m f^{(m)} > 0$ on the same interval the block sum is strictly positive (`Fabius.thueMorseBlockSum_pos`), by induction through the one-step reduction $\sum_{n < 2^{m+1}} \varepsilon_n f(x + nh) = - \int_0^h \sum_{n < 2^m} \varepsilon_n f'(x + t + 2nh) dt$ (`Fabius.thueMorseBlockSum_succ_eq_neg_integral`), whose integrand is a finite sum of continuous functions and so needs no parametric-integral theory for its integrability. What remains informal in the two applications (9.8) and (9.9) is only the computation of the explicit derivatives of $(a+y)^{-s}$ and $\log(a+y)$; their signs then feed `Fabius.thueMorseBlockSum_pos` directly.

9.3 Arbitrary translation weights

The dyadic steps $h, 2h, 4h, \dots$ give the identities above their Thue–Morse reading, but the cancellation mechanism does not use them. The formalization isolates the general statement, which is worth recording: it subsumes the ordinary and sparse Prouhet identities simultaneously and explains exactly where the powers of two matter (only in the value of the sharp moment, never in the vanishing).

Theorem 9.4 (Arbitrary-weight Prouhet identities). *Let R be a commutative ring, S a finite index set, $w: S \rightarrow R$ arbitrary weights, and $x \in R$. For every polynomial $p \in R[y]$ with $\deg p < |S|$,*

$$\sum_{T \subseteq S} (-1)^{|T|} p\left(x + \sum_{j \in T} w_j\right) = 0, \quad (9.10)$$

and at the first potentially surviving degree,

$$\sum_{T \subseteq S} (-1)^{|T|} \left(x + \sum_{j \in T} w_j\right)^{|S|} = (-1)^{|S|} |S|! \prod_{j \in S} w_j. \quad (9.11)$$

The word “potentially” is essential over a general commutative ring. Equation (9.11) is an unconditional identity, but its displayed factor can vanish: a step may be zero or a zero divisor, and the image of $|S|!$ may vanish in positive characteristic. Thus the formula identifies the exact top moment without asserting that the cancellation order is always exactly $|S|$. For the classical unit-step dyadic weights $w_j = 2^j$ over $\mathbb{Z}$, all of these factors are nonzero.

Proof. For (9.10), induct on S . If $S = \emptyset$, the strict degree condition forces $p = 0$, because the zero polynomial has degree $-\infty$; the sole powerset summand is therefore zero. For $S \cup \{a\}$ with $a \notin S$, split the subsets by membership of a and pair T with $T \cup \{a\}$:

$$\sum_{T \subseteq S \cup \{a\}} (-1)^{|T|} p(x + w_T) = \sum_{T \subseteq S} (-1)^{|T|} q(x + w_T), \quad q(y) := p(y) - p(y + w_a),$$

where $w_T := \sum_{j \in T} w_j$. The translation $p(y + w_a)$ has the same degree and the same leading coefficient as p , so $\deg q < \deg p$ (and $q = 0$ when p is constant). The induction hypothesis applied to q finishes.

For (9.11), induct on S again with $p(y) = y^{s+1}$, $s = |S|$. The same pairing produces $q(y) = y^{s+1} - (y + w_a)^{s+1} = -\sum_{k \leq s} \binom{s+1}{k} w_a^{s+1-k} y^k$ by the binomial theorem, the top term cancelling. Summing over $T \subseteq S$ and exchanging the two finite sums, every row $k < s$ vanishes by (9.10), and the row $k = s$ contributes $-\binom{s+1}{s} w_a \cdot (-1)^s s! \prod_{j \in S} w_j$ by the induction hypothesis. Since $\binom{s+1}{s} = s+1$, the total is $(-1)^{s+1} (s+1)! w_a \prod_{j \in S} w_j$, as claimed. $\square$

Formal versions. The degree-valued cancellation, including the zero polynomial on the empty support, is `Fabius.sum_powerset_neg_one_pow_eval_of_degree_lt`; its legacy natural-degree companion is `Fabius.sum_powerset_neg_one_pow_eval`. The sharp moment is `Fabius.sum_powerset_neg_one_pow_pow_card`, in `ThueMorseSparseProuhet.lean`; the degree drop is Mathlib’s `Polynomial.taylor` bookkeeping. The reusable passage from lower monomial moments to polynomial extraction is isolated in `FinitePolynomialFunctional.lean`: in particular `Fabius.sum_weight_mul_eval_eq_coeff_mul_of_moments` and `Fabius.sum_weight_mul_eval_eq_zero_of_degree_lt` are its same-ring extraction and cancellation interfaces. The module also supplies scalar-extension forms; none assumes distinct nodes or a nonzero top moment. The direct dyadic block wrappers obtained by specializing $w_j = 2^j h$ are recorded next.

9.4 Prouhet cancellation and the first surviving moment

Theorem 9.5 (Prouhet polynomial annihilation). *If p is a polynomial of degree less than m , then*

$$\sum_{n=0}^{2^m-1} \varepsilon_n p(x + nh) = 0. \quad (9.12)$$

If $p(y) = cy^m + \text{lower-degree terms}$, then

$$\sum_{n=0}^{2^m-1} \varepsilon_n p(x + nh) = (-1)^m m! 2^{\binom{m}{2}} h^m c. \quad (9.13)$$

Proof. Every m -fold finite difference annihilates a polynomial of degree below m . At degree exactly m , its m -th derivative is the constant $m!c$, so (9.5) gives the sharp value. $\square$

Theorem 9.6 (Commutative-ring dyadic block extraction). *Let R be a commutative ring, let $p \in R[y]$ have degree at most m , and let $x, h \in R$. With the integer signs mapped canonically into R , the whole affine block is the degree- m coefficient functional*

$$\sum_{n=0}^{2^m-1} \varepsilon_n p(x + nh) = [y^m] p \left((-1)^m (m! \cdot 1_R) 2^{\binom{m}{2}} h^m \right). \quad (9.14)$$

Consequently, with the degree of the zero polynomial understood as $-\infty$,

$$\deg p < m \implies \sum_{n=0}^{2^m-1} \varepsilon_n p(x + nh) = 0. \quad (9.15)$$

Both statements include $m = 0$: the extraction formula is evaluation of a constant polynomial, while the strict-degree hypothesis then admits only the zero polynomial.

Proof. Expand p in monomials and interchange the two finite sums. The moments below m vanish by (9.10), and (9.11) gives the degree- m moment for the steps $2^j h$. Thus only $[y^m]p$ remains. The cancellation statement is the same calculation with that coefficient equal to zero. $\square$

Formal versions. The extraction identity (9.14) is `Fabius.sum_thueMorseSign_mul_affine_eval_eq_coeff_card`; the degree-valued cancellation (9.15) is `Fabius.sum_thueMorseSign_mul_affine_eval_of_degree_lt`. Their monomial corollaries are `Fabius.sum_thueMorseSign_mul_affine_pow_eq_zero` and `Fabius.sum_thueMorseSign_mul_affine_pow_card`, all in `ThueMorseSparseProuhet.lean`. These are exact identities over every commutative ring. In positive characteristic the scalar $(m! \cdot 1_R)2^{\binom{m}{2}}h^m$ may be zero, so the formal theorem deliberately does not claim that a nonzero top coefficient produces a nonzero block sum.

For $h = 1$, let

$$A_m = \{0 \leq n < 2^m : \tau(n) = 0\}, \quad B_m = \{0 \leq n < 2^m : \tau(n) = 1\}. \quad (9.16)$$

Then

$$\sum_{n \in A_m} n^r = \sum_{n \in B_m} n^r, \quad 0 \leq r < m, \quad (9.17)$$

which is Prouhet’s classical equal-sums-of-like-powers construction. Over $\mathbb{Z}$, at $r = m$ the difference is nonzero, so the cancellation order of this classical unit-step partition is exact.

Remark 9.7 (The Prouhet–Tarry–Escott problem). Equation (9.17) is the historical origin of the *Prouhet–Tarry–Escott problem*: given a degree d , find two disjoint multisets $A \neq B$ of integers of equal size with $\sum_{a \in A} a^r = \sum_{b \in B} b^r$ for all $0 \leq r \leq d$. A solution is called *ideal* when its size is $d + 1$, the smallest conceivable. Prouhet’s partition (9.16) solves degree $d = m - 1$ with the much larger size 2^{m-1} , but it does so uniformly for every degree – still the only completely elementary construction that scales [2, 17, 18]. Ideal solutions are known for all $d \leq 9$ and for $d = 11$; their existence for every degree is open. The sparse Prouhet theorem of section 10 produces further nonconsecutive solutions: for any n with $\text{wt}_2(n) = s$, splitting the submasks of n by sign parity gives two sets of size 2^{s-1} with equal power sums through degree $s - 1$, supported on the odd positions of Pascal row n rather than on an interval.

9.5 The finite exponential transform

Define

$$M_{m,r} = \sum_{n=0}^{2^m-1} \varepsilon_n n^r. \quad (9.18)$$

Theorem 9.8 (Finite complex Hamming-weight and Thue–Morse transforms). *For every $m \in \mathbb{N}_0$ and $u, a, z \in \mathbb{C}$,*

$$\sum_{n < 2^m} u^{\text{wt}_2(n)} e^{nz} = \prod_{j < m} (1 + u e^{2^j z}), \quad (9.19)$$

$$\sum_{n < 2^m} u^{\text{wt}_2(n)} e^{a+nz} = e^a \prod_{j < m} (1 + u e^{2^j z}), \quad (9.20)$$

$$\sum_{n < 2^m} \varepsilon_n e^{nz} = \prod_{j < m} (1 - e^{2^j z}), \quad (9.21)$$

$$\sum_{n < 2^m} \varepsilon_n e^{a+nz} = e^a \prod_{j < m} (1 - e^{2^j z}). \quad (9.22)$$

All four identities include $m = 0$, with the usual empty-product convention. No convergence hypothesis is present: every sum and product is finite.

Proof. Expand the product in (9.19). Choosing the $ue^{2^j z}$ term at the bit positions of $n < 2^m$ contributes $u^{\text{wt}_2(n)}e^{nz}$, because the selected powers of two sum to n . The exponential addition law factors the common translate e^a and gives (9.20). Finally put $u = -1$ and use $\varepsilon_n = \varepsilon_n$ to obtain (9.21) and (9.22). $\square$

Formal versions. The four displayed identities are, in order, `Fabius.sum_range_two_pow_binaryWeight_cexp`, `Fabius.sum_range_two_pow_binaryWeight_cexp_affine`, `Fabius.sum_range_two_pow_thueMorseSign_cexp`, and `Fabius.sum_range_two_pow_thueMorseSign_cexp_affine` in `ThueMorseComplexExponential.lean`. Their parameters are genuinely complex, their affine shifts are arbitrary, and their statements include the empty block. They are finite identities only: this module asserts neither an infinite product nor any sine or sinc normalization.

Taking $z = t$ in (9.21) and expanding the finite sum of exponentials gives

$$\sum_{r \geq 0} M_{m,r} \frac{t^r}{r!} = \sum_{n < 2^m} \varepsilon_n e^{nt} = \prod_{j=0}^{m-1} (1 - e^{2^j t}). \quad (9.23)$$

The order- m zero at $t = 0$ immediately gives

$$M_{m,r} = 0 \quad (r < m), \quad M_{m,m} = (-1)^m m! 2^{\binom{m}{2}}. \quad (9.24)$$

Further formal versions. The vanishing and sharp value (9.24) have long been formal — for an arbitrary rational translation — as `Fabius.thueMorseTranslatedPowerSum_eq_zero_of_lt` and `Fabius.thueMorseTranslatedPowerSum_self` (`ThueMorseExponential.lean`); the raw moments $M_{m,r}$ are the translation $c = 2^m$ of that interface, `Fabius.thueMorseTranslatedPowerSum_two_pow` (`ThueMorseMoments.lean`). In (9.23), the equality between the finite exponential sum and the finite product is now exactly `Fabius.sum_range_two_pow_thueMorseSign_cexp`, for every complex t . The first equality, which reads that finite exponential sum as an analytic exponential generating function, is formal at the coefficient level: the dictionary $[t^r] \sum_{n < 2^m} \varepsilon_n e^{nt} = M_{m,r}/r!$ is `Fabius.coeff_sum_thueMorseSign_exp_pow`, and the right-hand product is the algebraic master product evaluated at $z = e^t \in \mathbb{Q}[[t]]$. No named Lean theorem in this crosswalk identifies that coefficientwise formal series with an analytically convergent EGF; the new complex theorem delivers precisely the finite middle-to-right identity, without silently adding that separate analytic bridge.

The normalized translated series itself has the exact finite product

$$\begin{aligned} \mathcal{S}_{c,k}(X) &:= \sum_{n \geq 0} \frac{\sum_{r < 2^k} \varepsilon_r (r - 2^k + c)^{n+k}}{(n+k)!} X^n, \\ \mathcal{S}_{c,k}(X) &= (-1)^k 2^{\binom{k}{2}} e^{(c-1)X} \prod_{j < k} \frac{e^{-2^j X} - 1}{-2^j X} \quad (c \in \mathbb{Q}, k \in \mathbb{N}_0). \end{aligned} \quad (9.25)$$

This is an identity in $\mathbb{Q}[[X]]$, where each displayed quotient denotes the removable formal series with constant coefficient one. It is `Fabius.thueMorseTranslatedShiftedPowerSeries_eq_exp1_product`; the quotient series is `Fabius.rationalExp1DivSeries`, both in `ThueMorseExponential.lean`.

9.6 All power moments: a finite composition formula

Theorem 9.9 (Complete moment composition). *For $m \geq 1$ and $r \geq m$,*

$$M_{m,r} = (-1)^m r! 2^{\binom{m}{2}} \sum_{\substack{q_0, \dots, q_{m-1} \geq 0 \\ q_0 + \dots + q_{m-1} = r-m}} \prod_{j=0}^{m-1} \frac{2^{jq_j}}{(q_j + 1)!}. \quad (9.26)$$

Equivalently,

$$M_{m,r} = (-1)^m r! \sum_{\substack{p_0, \dots, p_{m-1} \geq 1 \\ p_0 + \dots + p_{m-1} = r}} \frac{2^{\sum_{j=0}^{m-1} jp_j}}{p_0! \cdots p_{m-1}!}. \quad (9.27)$$

Proof. Factor the lowest-order term from each factor in (9.23):

$$\begin{aligned} \prod_{j=0}^{m-1} (1 - e^{2^j t}) &= (-1)^m 2^{\binom{m}{2}} t^m \prod_{j=0}^{m-1} \frac{e^{2^j t} - 1}{2^j t} \\ &= (-1)^m 2^{\binom{m}{2}} t^m \prod_{j=0}^{m-1} \sum_{q \geq 0} \frac{2^{jq} t^q}{(q+1)!}. \end{aligned}$$

Extract t^r and multiply by $r!$. Set $p_j = q_j + 1$ for the second form. $\square$

Formal versions. `ThueMorseMoments.lean` proves [theorem 9.9](#) as `Fabius.sum_thueMorseSign_mul_pow_add`: for all $d \geq 0$,

$$\sum_{n < 2^m} \varepsilon_n n^{m+d} = (-1)^m (m+d)! \sum_{\substack{q_0, \dots, q_{m-1} \geq 0 \\ q_0 + \dots + q_{m-1} = d}} \prod_{j=0}^{m-1} \frac{(2^j)^{q_j+1}}{(q_j + 1)!},$$

with the compositions realized as finitely supported functions (`Finset.finsuppAntidiag`). Factoring $(2^j)^{q_j+1} = 2^j \cdot 2^{jq_j}$ across the product recovers the boxed constant $2^{\binom{m}{2}}$ of (9.26), and $p_j = q_j + 1$ gives (9.27). The substitution $z = e^t$ is performed once and for all on an *arbitrary* support of bit positions – each factor becomes $-t$ times an explicit entire series (`Fabius.one_sub_exp_pow`, specialized in `Fabius.one_sub_exp_two_pow`, and collected in `Fabius.prod_one_sub_exp_two_pow`) – and the $t^{|S|+d}$ coefficient is extracted through `PowerSeries.coeff_prod` inside `Fabius.sum_powerset_thueMorseSign_mul_pow_add`. The block statement is then its case $S = \{0, \dots, m-1\}$, transported across the Boolean-cube reindexing kernel `Fabius.sum_powerset_two_pow` from `ThueMorseBitSupport.lean`.

9.7 Bell-polynomial compression and cumulants

Let $B_{2\ell}$ be the Bernoulli numbers with $B_2 = 1/6$, and define

$$\kappa_1^{(m)} = \frac{2^m - 1}{2}, \quad (9.28)$$

$$\kappa_{2\ell}^{(m)} = \frac{B_{2\ell}}{2\ell} \frac{2^{2\ell m} - 1}{2^{2\ell} - 1} \quad (\ell \geq 1), \quad (9.29)$$

$$\kappa_{2\ell+1}^{(m)} = 0 \quad (\ell \geq 1). \quad (9.30)$$

Theorem 9.10 (Bell-polynomial moments). *For $m \geq 1$ and $r \geq 0$,*

$$M_{m,m+r} = (-1)^m 2^{\binom{m}{2}} \frac{(m+r)!}{r!} \mathcal{B}_r^{\exp} \left(\kappa_1^{(m)}, \dots, \kappa_r^{(m)} \right). \quad (9.31)$$

Proof. The classical formal expansion

$$\log\left(\frac{e^y - 1}{y}\right) = \frac{y}{2} + \sum_{\ell \geq 1} \frac{B_{2\ell}}{2\ell(2\ell)!} y^{2\ell} \quad (9.32)$$

turns the normalized product into

$$\prod_{j=0}^{m-1} (1 - e^{2^j t}) = (-1)^m 2^{\binom{m}{2}} t^m \exp\left(\sum_{r \geq 1} \kappa_r^{(m)} \frac{t^r}{r!}\right). \quad (9.33)$$

Use the defining generating function (7.6) and compare coefficients. $\square$

The first three nonzero layers are

$$M_{m,m} = (-1)^m 2^{\binom{m}{2}} m!, \quad (9.34)$$

$$M_{m,m+1} = (-1)^m 2^{\binom{m}{2}} (m+1)! \frac{2^m - 1}{2}, \quad (9.35)$$

$$M_{m,m+2} = (-1)^m 2^{\binom{m}{2}} (m+2)! \left(\frac{(2^m - 1)^2}{8} + \frac{4^m - 1}{72} \right). \quad (9.36)$$

9.8 Centered moments and a parity selection rule

Let

$$c_m = \frac{2^m - 1}{2}. \quad (9.37)$$

The centered exponential transform is especially symmetric.

Theorem 9.11 (Centered sinh product). *For every $m \geq 0$,*

$$\boxed{\sum_{n=0}^{2^m-1} \varepsilon_n e^{(n-c_m)t} = (-2)^m \prod_{j=0}^{m-1} \sinh(2^{j-1}t).} \quad (9.38)$$

Consequently

$$\boxed{\sum_{n < 2^m} \varepsilon_n (n - c_m)^r = 0} \quad (9.39)$$

whenever $r < m$ or $r \not\equiv m \pmod{2}$. At $r = m$,

$$\sum_{n < 2^m} \varepsilon_n (n - c_m)^m = (-1)^m m! 2^{\binom{m}{2}}. \quad (9.40)$$

Proof. Use $1 - e^u = -2e^{u/2} \sinh(u/2)$ in (9.23). The product of exponential factors is $e^{c_m t}$, which disappears after centering. The right side is a product of m odd functions, hence has parity m , and each factor starts with a linear term. Coefficient comparison proves the selection rule and the first coefficient. $\square$

Remark 9.12 (Reflection proof of the selection rule). The parity cancellation (9.39) needs no exponential transform. Reindex $n \mapsto 2^m - 1 - n$ in $S_r = \sum_{n < 2^m} \varepsilon_n (n - c_m)^r$: the reflection formula (2.8) gives $\varepsilon_{2^m-1-n} = (-1)^m \varepsilon_n$, while $(2^m - 1 - n) - c_m = -(n - c_m)$. Hence $S_r = (-1)^{m+r} S_r$, and $S_r = 0$ whenever $m + r$ is odd. The same reindexing at an *arbitrary* translation x yields the functional equation

$$\sum_{n < 2^m} \varepsilon_n (x + n)^r = (-1)^{m+r} \sum_{n < 2^m} \varepsilon_n (n - x - (2^m - 1))^r,$$

valid over any commutative ring; the selection rule is its fixed point at the self-reflective translation $x = -c_m$.

Formal boundary. At the raw exponential level, the centered precursor

$$\sum_{n < 2^m} \varepsilon_n e^{(n - c_m)t} = e^{-c_m t} \prod_{j < m} (1 - e^{2^j t})$$

is formal for every complex t : it is `Fabius.sum_range_two_pow_thueMorseSign_cexp_affine` with $a = -c_m t$ and $z = t$. The subsequent complex half-angle rewrite $1 - e^u = -2e^{u/2} \sinh(u/2)$, and hence the centered sinh product (9.38) as a named complex theorem, is not yet packaged in Lean. At the moment level, the reflection functional equation is `Fabius.sum_thueMorseSign_mul_add_pow_reflect` (`ThueMorseMoments.lean`, any commutative ring); the selection rule (9.39) for odd $m + r$ is `Fabius.sum_thueMorseSign_mul_midpoint_pow_eq_zero`, proved by the reflection argument of [remark 9.12](#) rather than through (9.38); the complementary case $r < m$ of either parity is `Fabius.thueMorseTranslatedPowerSum_eq_zero_of_lt`. The first coefficient (9.40) is `Fabius.sum_thueMorseSign_mul_midpoint_pow_self`.

This extra parity cancellation is invisible in the uncentered Prouhet statement. It is a consequence of combining dyadic reflection with moment cancellation.

9.9 All binomial-basis moments

Define

$$B_{m,r} = \sum_{n=0}^{2^m-1} \varepsilon_n \binom{n}{r}. \quad (9.41)$$

Theorem 9.13 (Binomial moment product). *For $\ell \geq 0$,*

$$B_{m,m+\ell} = (-1)^m 2^{\binom{m}{2}} [u^\ell] \prod_{j=0}^{m-1} \frac{(1+u)^{2^j} - 1}{2^j u}. \quad (9.42)$$

For $r < m$, $B_{m,r} = 0$; in particular,

$$B_{m,m} = (-1)^m 2^{\binom{m}{2}}. \quad (9.43)$$

Proof. Expand $P_m(1+u)$ in two ways:

$$P_m(1+u) = \sum_{n < 2^m} \varepsilon_n (1+u)^n = \sum_{r \geq 0} B_{m,r} u^r,$$

while

$$P_m(1+u) = (-1)^m 2^{\binom{m}{2}} u^m \prod_{j=0}^{m-1} \frac{(1+u)^{2^j} - 1}{2^j u}.$$

Compare coefficients. □

Corollary 9.14 (Integer composition form). *For all $m, \ell \geq 0$,*

$$B_{m,m+\ell} = (-1)^m \sum_{\substack{q_0, \dots, q_{m-1} \geq 0 \\ q_0 + \dots + q_{m-1} = \ell}} \prod_{j=0}^{m-1} \binom{2^j}{q_j + 1}. \quad (9.44)$$

Proof. In (9.42) each normalized factor is, up to its constant 2^j , the geometric sum $((1+u)^{2^j} - 1)/u = \sum_{i < 2^j} (1+u)^i$, whose q -th coefficient is $\sum_{i < 2^j} \binom{i}{q} = \binom{2^j}{q+1}$ by the hockey-stick identity. Multiplying the m factors and collecting the coefficient of u^ℓ gives (9.44). The single composition at $\ell = 0$ recovers (9.43), since $\binom{2^j}{1} = 2^j$ multiplies to $2^{\binom{m}{2}}$. Unlike the power-moment composition (9.26), every summand here is an integer. □

Formal versions. `ThueMorseBlockAlgebra.lean` proves the substitution identity behind (9.42) as `Fabius.sum_thueMorseSign_mul_one_add_X_pow` (over any commutative ring), the vanishing $B_{m,r} = 0$ for $r < m$ as `Fabius.sum_thueMorseSign_mul_choose_eq_zero`, and the sharp value (9.43) as `Fabius.sum_thueMorseSign_mul_choose_self`, each with an integer specialization. The complete composition form (9.44) is `Fabius.sum_thueMorseSign_mul_choose_add` (`ThueMorseMoments.lean`), with the hockey-stick column sum $\sum_{i < N} \binom{i}{q} = \binom{N}{q+1}$ isolated as the reusable `Fabius.sum_range_choose_eq_choose_succ`.

The binomial basis is often preferable in discrete formalization because finite differences obey the simple lowering rule $\Delta \binom{x}{r} = \binom{x}{r-1}$.

10 Sparse Prouhet identities on arbitrary Pascal supports

The ordinary Prouhet block $[0, 2^m)$ is the submask lattice of $2^m - 1$. Lucas’s theorem suggests replacing this all-ones mask by an arbitrary n .

Define

$$J(n) = \{j \geq 0 : b_j(n) = 1\}, \quad s = |J(n)| = \text{wt}_2(n), \quad \beta(n) = \sum_{j \in J(n)} j. \quad (10.1)$$

Theorem 10.1 (Signed Pascal-support product). *For every $n \geq 0$,*

$$R_n(z) := \sum_{k \preceq_2 n} \varepsilon_k z^k = \prod_{j \in J(n)} (1 - z^{2^j}). \quad (10.2)$$

Equivalently,

$$R_n(z) = \sum_{k=0}^n \varepsilon_k \mathbf{1}_{\binom{n}{k} \text{ odd}} z^k. \quad (10.3)$$

Proof. Expanding the product chooses a subset of the occupied bit positions of n . Its exponent is the corresponding submask k , and its sign is $\varepsilon_k = \varepsilon_k$. Lucas’s criterion gives the second form. $\square$

Corollary 10.2 (Balance on odd Pascal positions). *For every $n > 0$,*

$$\sum_{\substack{0 \leq k \leq n \\ \binom{n}{k} \text{ odd}}} \varepsilon_k = 0. \quad (10.4)$$

Thus exactly half of the odd positions in every nonzero Pascal row have even binary weight and half have odd binary weight.

Proof. Set $z = 1$ in (10.2). There are $2^{\text{wt}_2(n)}$ support positions by [corollary 5.2](#). $\square$

Theorem 10.3 (Sparse Prouhet theorem). *Let $n \geq 0$, $s = \text{wt}_2(n)$. For every polynomial p of degree below s ,*

$$\sum_{k \preceq_2 n} \varepsilon_k p(k) = 0. \quad (10.5)$$

The first surviving power moment is

$$\sum_{k \preceq_2 n} \varepsilon_k k^s = (-1)^s s! 2^{\beta(n)}. \quad (10.6)$$

Proof. Every factor in (10.2) has a simple zero at $z = 1$, so R_n has a zero of exact order s , with local leading term

$$R_n(z) = (-1)^s 2^{\beta(n)} (z - 1)^s + \mathcal{O}((z - 1)^{s+1}).$$

The derivative/falling-factorial proof of [theorem 9.5](#) applies with the sum restricted to submasks. When $n = 0$, one has $s = 0$: the cancellation hypothesis forces $p = 0$, while the sharp moment reads $0^0 = 1$ on the left and 1 on the right under the usual polynomial convention. $\square$

Formal versions. The canonical bit support `Fabius.bitSupport` and its dictionary `Fabius.mem_bitSupport`, `Fabius.sum_two_pow_bitSupport`, `Fabius.card_bitSupport` are in `ThueMorseBitSupport.lean`; its global `Fabius.sum_two_pow_injective` makes the powerset parametrization faithful. In `ThueMorseSparseProuhet.lean`, the sparse cancellation and the sharp moment are respectively the all-index degree-valued theorem `Fabius.sum_submask_neg_one_pow_eval_of_degree_lt` and `Fabius.sum_submask_neg_one_pow_pow`, both instances of the arbitrary-weight engine of [theorem 9.4](#) with $w_j = 2^j h$ running over $J(n)$, over any commutative ring. The generating identity (10.2) itself is now `Fabius.prod_one_sub_pow_bitSupport` (`ThueMorseLucasSupport.lean`), the submasks running over the powerset of the bit support; through Lucas’s criterion this index set is exactly `Fabius.oddBinomialIndices`, and the literal odd-Pascal-index forms of the two displayed identities are `Fabius.sum_oddBinomialIndices_thueMorseSign_mul_affine_eval_of_degree_lt` and `Fabius.sum_oddBinomialIndices_thueMorseSign_mul_affine_pow_binaryWeight`. The balance [corollary 10.2](#) is `Fabius.sum_thueMorseSign_oddBinomialIndices`.

Corollary 10.4 (Sparse operator identity). *For every function f ,*

$$\sum_{k \preceq_{2^n} n} \varepsilon_k f(x + hk) = \prod_{j \in J(n)} (I - T_{2^j h}) f(x). \quad (10.7)$$

If $f \in C^s$ and $h \geq 0$, its absolute value is at most

$$2^{\beta(n)} h^s \sup_{x \leq y \leq x + nh} |f^{(s)}(y)|. \quad (10.8)$$

Proof. Expand the product over the set-bit positions and apply the fundamental theorem of calculus once for each step. $\square$

The whole sparse moment sequence also has a direct finite formula.

Corollary 10.5 (All sparse moments). *For $r \geq s$,*

$$\sum_{k \preceq_{2^n} n} \varepsilon_k k^r = (-1)^s r! 2^{\beta(n)} \sum_{\substack{(q_j)_{j \in J(n)} \geq 0 \\ \sum_{j \in J(n)} q_j = r - s}} \prod_{j \in J(n)} \frac{2^{jq_j}}{(q_j + 1)!}. \quad (10.9)$$

Proof. Substitute $z = e^t$ in (10.2), factor $-2^j t$ from every occupied bit, and extract coefficients exactly as in [theorem 9.9](#). $\square$

Formal versions. `ThueMorseSparseMoments.lean` proves [corollary 10.5](#) as `Fabius.sum_submask_thueMorseSign_mul_pow_add`, and that is now all it contains: the corollary is the special case $S = J(n)$ of a statement on an *arbitrary* finite set S of bit positions (`Fabius.sum_powerset_t_thueMorseSign_mul_pow_add`), which also subsumes the dyadic block $S = \{0, \dots, m - 1\}$ of [theorem 9.9](#). The arbitrary-support moment statement and the master product live one level down, in `ThueMorseMoments.lean` — an ancestor module cannot import a descendant, and [theorem 9.9](#) is itself derived from them. The master product on an arbitrary support is `Fabius.prod_one_sub_po`

`w_powerset`, over any commutative ring. Its hypothesis-free weight and sign inputs are centralized in `ThueMorseBitSupport.lean` as `Fabius.binaryWeight_sum_two_pow_eq_card` and its sign companion `Fabius.thueMorseSign_sum_two_pow`.

Example 10.6. For $n = 13 = (1101)_2$, $J(13) = \{0, 2, 3\}$, $s = 3$, and $\beta(13) = 5$. The odd positions in row 13 are

$$0, 1, 4, 5, 8, 9, 12, 13.$$

Their signed sums against $1, k, k^2$ vanish, while

$$\sum_{k \preceq_2 13} \varepsilon_k k^3 = (-1)^3 3! 2^5 = -192.$$

This is a Prouhet partition on a sparse, Pascal-determined grid rather than a consecutive interval.

11 Exact prefix sums, iterated antiderivatives, and the Fabius bridge

11.1 The first prefix discrepancy

Define the exclusive signed prefix and the number of one-bits in the sequence prefix by

$$S(N) = \sum_{n=0}^{N-1} \varepsilon_n, \quad A(N) = \sum_{n=0}^{N-1} \tau(n). \quad (11.1)$$

Theorem 11.1 (Exact prefix formula). *For every $q \geq 0$,*

$$\boxed{S(2q) = 0, \quad S(2q+1) = \varepsilon_q.} \quad (11.2)$$

Equivalently,

$$S(N) = \frac{1 - (-1)^N}{2} \varepsilon_{\lfloor N/2 \rfloor}. \quad (11.3)$$

Therefore

$$\boxed{A(N) = \frac{N - S(N)}{2}} \quad (11.4)$$

and explicitly

$$A(2q) = q, \quad A(2q+1) = q + \tau(q). \quad (11.5)$$

Proof. Every adjacent pair cancels: $\varepsilon_{2j} + \varepsilon_{2j+1} = \varepsilon_j - \varepsilon_j = 0$. An odd-length prefix leaves the unpaired term $\varepsilon_{2q} = \varepsilon_q$. The count formula follows from $\varepsilon_n = 1 - 2\tau(n)$. $\square$

The inclusive sum through N is

$$\sum_{n=0}^N \varepsilon_n = \frac{1 + (-1)^N}{2} \varepsilon_{\lfloor N/2 \rfloor}. \quad (11.6)$$

Corollary 11.2 (Interval discrepancy). *For $0 \leq a \leq b$,*

$$\sum_{n=a}^{b-1} \varepsilon_n = S(b) - S(a), \quad (11.7)$$

$$\left| \sum_{n=a}^{b-1} \varepsilon_n \right| \leq 2, \quad (11.8)$$

$$\left| \sum_{n=a}^{b-1} \tau(n) - \frac{b-a}{2} \right| \leq 1. \quad (11.9)$$

Both constants are attained.

Proof. The interval sum is a difference of two values in $\{-1, 0, 1\}$. Convert to zero–one values by (1.3). For sharpness, $\varepsilon_5 + \varepsilon_6 = 2$. $\square$

Formal versions. `ThueMorseEulerTransform.lean` proves the exclusive closed form (11.2) as `Fabius.sum_thueMorseSign_range` (even case `Fabius.sum_thueMorseSign_range_two_mul`), the inclusive form (11.6) as `Fabius.sum_thueMorseSign_range_succ`, the ones-count formulas (11.5) as `Fabius.sum_thueMorseBit_range_two_mul` and `Fabius.sum_thueMorseBit_range_two_mul_add_one`, and the bridge $2A(N) + S(N) = N$ as `Fabius.two_mul_sum_thueMorseBit_add_sum_thueMorseSign`. The three displayed statements of corollary 11.2 are `Fabius.sum_thueMorseSign_Ico` for (11.7), `Fabius.abs_sum_thueMorseSign_Ico_le_two` for (11.8) – both resting on the prefix bound $|S(N)| \leq 1$ (`Fabius.abs_sum_thueMorseSign_range_le_one`) – and, for the bits in doubled integer form, `Fabius.abs_two_mul_sum_thueMorseBit_Ico_sub_le_two` for (11.9). All four are inequalities or exact identities; the attainment of the two constants, checked in the proof above at $a = 5, b = 7$, is not formalized.

Corollary 11.3 (Aligned and partial dyadic blocks). *For $m \geq 1$ and $q \geq 0$,*

$$\sum_{r=0}^{2^m-1} \varepsilon_{2^m q+r} = 0. \quad (11.10)$$

More generally, for $0 \leq R \leq 2^m$,

$$\sum_{r=0}^{R-1} \varepsilon_{2^m q+r} = \varepsilon_q S(R). \quad (11.11)$$

Proof. Use (2.6) and factor out ε_q ; then apply theorem 11.1. $\square$

Formal versions. (11.11) and (11.10) are `Fabius.sum_thueMorseSign_block_prefix` and `Fabius.sum_thueMorseSign_aligned_block` (`ThueMorseEulerTransform.lean`).

The prefix generating function is unusually simple:

$$\sum_{N \geq 0} S(N) z^N = \frac{z}{1-z} E(z) = z E(z^2) = z \prod_{j \geq 1} (1 - z^{2^j}). \quad (11.12)$$

The middle equality is the Mahler equation.

Formal versions. The first and last members of (11.12) are formal: `Fabius.thueMorsePrefixSeries_eq` (`ThueMorseMahler.lean`) proves $\sum_N S(N) z^N = z E(z^2)$ directly from the exact prefix collapse, which under the dissection is ‘ $z \prod_{j \geq 1} (1 - z^{2^j})$ ’ at every truncation level.

11.2 Iterated inclusive prefixes

Define discrete antiderivatives by

$$S_n^{(0)} = \varepsilon_n, \quad S_n^{(k+1)} = \sum_{j=0}^n S_j^{(k)}. \quad (11.13)$$

Then (11.6) says

$$S_{2q}^{(1)} = \varepsilon_q, \quad S_{2q+1}^{(1)} = 0. \quad (11.14)$$

Theorem 11.4 (Binomial convolution for discrete antiderivatives). *For $k \geq 1$,*

$$S_n^{(k)} = \sum_{j=0}^n \binom{n-j+k-1}{k-1} \varepsilon_j. \quad (11.15)$$

Its ordinary generating function is

$$\sum_{n \geq 0} S_n^{(k)} z^n = \frac{E(z)}{(1-z)^k}. \quad (11.16)$$

Proof. One inclusive summation multiplies an ordinary generating function by $(1-z)^{-1}$. Iterate k times, then expand $(1-z)^{-k} = \sum_{r \geq 0} \binom{r+k-1}{k-1} z^r$ and take the Cauchy product. $\square$

Formal versions. The corpus module `ThueMorsePrefix.lean` defines `Fabius.iteratedPrefix` and proves the binomial convolution (11.15) as `Fabius.iteratedPrefix_convolution`; the generating identity (11.16) is formal below every dyadic truncation as `Fabius.coeff_thueMorseBlockPolynomial_mul_invOneSubPow_eq_iteratedPrefix` (`ThueMorseGenerating.lean`).

11.3 Terminal zero runs at dyadic boundaries

The quotient of a finite Prouhet polynomial by $(1-z)^k$ is still a polynomial when $k \leq r$:

$$\frac{P_r(z)}{(1-z)^k} = (1-z)^{r-k} \prod_{j=0}^{r-1} (1+z+\dots+z^{2^j-1}). \quad (11.17)$$

Its degree is $2^r - k - 1$.

Theorem 11.5 (Dyadic terminal zero run). *If $1 \leq k \leq r$, then*

$$S_{2^r-k+i}^{(k)} = 0, \quad 0 \leq i < k. \quad (11.18)$$

The adjacent values are

$$S_{2^r-k-1}^{(k)} = (-1)^{r-k}, \quad S_{2^r}^{(k)} = -1. \quad (11.19)$$

Proof. Below degree 2^r , the coefficients of $E(z)/(1-z)^k$ agree with those of $P_r(z)/(1-z)^k$. The latter has degree $2^r - k - 1$, proving the zero run. Its leading coefficient is $(-1)^{r-k}$, giving the preceding value. The coefficient at 2^r comes from the first omitted factor $1 - z^{2^r}$, or directly from the prefix recurrence. $\square$

Formal versions. All three statements of [theorem 11.5](#) have long been formal in `ThueMorsePrefix.lean`: `Fabius.iteratedPrefix_dyadic_zero_run` (the run), `Fabius.iteratedPrefix_before_dyadic_run` ($S_{2^r-k-1}^{(k)} = (-1)^{r-k}$), and `Fabius.iteratedPrefix_at_dyadic` ($S_{2^r}^{(k)} = -1$), with the endpoint value `Fabius.iteratedPrefix_dyadic_endpoint`. The factorization behind (11.17) is the geometric ladder `Fabius.prod_one_sub_two_pow_eq_geom` and `Fabius.thueMorseBlockPolynomial_eq_geom_prod` (`ThueMorseBlockAlgebra.lean`), over any commutative ring.

Corollary 11.6 (Signed reversal before the terminal run). *If $0 \leq k \leq r$ and $k + d < 2^r$, then the admissible coefficient window before the final k zeros is reciprocal up to the parity sign $(-1)^{r-k}$:*

$$S_{2^r-k-1-d}^{(k)} = (-1)^{r-k} S_d^{(k)}. \quad (11.20)$$

In particular, reversal preserves its zero locus exactly:

$$S_{2^r-k-1-d}^{(k)} = 0 \iff S_d^{(k)} = 0. \quad (11.21)$$

The case $k = 0$ is the usual dyadic complement law for the signs, and $d = 0$ recovers the value immediately preceding the terminal run.

Proof. Compare opposite coefficients in the reciprocal form of (11.17). The only multiplier is $(-1)^{r-k}$, which is a unit in $\mathbb{Z}$; hence multiplication by it preserves and reflects vanishing. $\square$

Formal versions. The signed identity (11.20) is `Fabius.iteratedPrefix_dyadic_reverse_window`, and its exact zero-locus consequence (11.21) is `Fabius.iteratedPrefix_dyadic_reverse_window_eq_zero_iff` in the corpus module `ThueMorsePrefix.lean`. Both statements retain the boundary cases $k = 0$ and $d = 0$.

11.4 Why this is the discrete spine of the Fabius theory

The chain connecting Thue–Morse signs to the Fabius and Rvachev functions is now algebraically visible:

$$\begin{aligned} \text{binary parity character} &\implies P_r(z) = \prod_{j < r} (1 - z^{2^j}) \\ &\implies \frac{P_r(z)}{(1 - z)^k} \implies \text{iterated-prefix coefficient grids} \end{aligned} \tag{11.22}$$

The zero of order k at $z = 1$ creates the exact terminal flatness in [theorem 11.5](#); after the normalizations developed in the main repository exposition, the resulting coefficient grids and polygonal interpolants converge to the Fabius function. On the Fourier side, replacing z by a point on the unit circle turns the same finite product into a finite product of sine factors, whose scaled limit is the sinc product associated with the Rvachev up-function.

Thus the discrete-prefix and analytic-sinc appearances of Thue–Morse are not separate phenomena. They are the coefficient and Fourier-evaluation faces of the single factorization (6.3). The quantitative convergence theory belongs to the dedicated companion reports and Lean modules; the purpose here is to expose the exact common algebraic spine.

12 Ordinary Fourier transforms and product-to-sum identities

12.1 A finite Fourier product

Putting $z = e^{i\theta}$ in (6.3) gives

$$\sum_{n=0}^{2^m-1} \varepsilon_n e^{in\theta} = \prod_{j=0}^{m-1} (1 - e^{i2^j\theta}). \tag{12.1}$$

Using $1 - e^{iu} = -2ie^{iu/2} \sin(u/2)$,

$$\sum_{n=0}^{2^m-1} \varepsilon_n e^{in\theta} = (-2i)^m e^{i(2^m-1)\theta/2} \prod_{j=0}^{m-1} \sin(2^{j-1}\theta). \tag{12.2}$$

Hence

$$\left| \sum_{n=0}^{2^m-1} \varepsilon_n e^{in\theta} \right| = 2^m \prod_{j=0}^{m-1} |\sin(2^{j-1}\theta)|. \tag{12.3}$$

Formal versions. The raw identity (12.1) is the specialization $z = i\theta$ of `Fabius.sum_range_two_pow_thueMorseSign_cexp` (`ThueMorseComplexExponential.lean`); it is valid for every *complex* θ , with no reality or convergence hypothesis. The affine companion `Fabius.sum_range_two_pow_thueMorseSign_cexp_affine` also allows an arbitrary complex translate of the exponent. Equivalently, the raw identity follows by evaluating the ring-valued master product `Fabius.prod_one_sub_pow_eq_sum_thueMorseSign` at $e^{i\theta}$. The sine form (12.2) is `Fabius.sum_thueMorseSign_exp_eq_sin_prod` (`ThueMorseSineProduct.lean`), built on

the half-angle factorization $1 - e^{iu} = -2ie^{iu/2} \sin(u/2)$ (Fabius.one_sub_exp_mul_I, proved by pure exponential algebra). The magnitude (12.3) is carried in squared form by the Riesz factorization `Fabius.normSq_sum_thueMorseSign_exp` (`ThueMorseFourier.lean`), whose factors are $|1 - e^{i\theta}|^2 = 2 - 2 \cos \theta$ (`Fabius.normSq_one_sub_exp`).

Theorem 12.1 (Affine real sine form). *For every $m \in \mathbb{N}_0$ and $x, \phi \in \mathbb{R}$,*

$$\sum_{n=0}^{2^m-1} \varepsilon_n \sin(\phi + nx) = 2^m \left(\prod_{j=0}^{m-1} \sin \frac{2^j x}{2} \right) \sin \left(\phi + \frac{(2^m - 1)x}{2} - \frac{m\pi}{2} \right). \quad (12.4)$$

The product is empty at $m = 0$.

Formal version. Equation (12.4) is exactly `Fabius.sum_thueMorseSign_mul_sin_affine` in `ThueMorseSineProduct.lean`. It follows from the complex sine-product form by an arbitrary phase rotation and extraction of imaginary parts.

12.2 Sine and cosine products as Thue–Morse sums

The affine identity specializes to the following exact real lacunary product: take $x = 2w$ and $\phi = m\pi/2 + w$, then append the resulting $\sin(2^m w)$ to the product over $j < m$.

Theorem 12.2 (Real lacunary sine product). *For every $m \in \mathbb{N}_0$ and $w \in \mathbb{R}$,*

$$2^m \prod_{j=0}^m \sin(2^j w) = \sum_{n=0}^{2^m-1} \varepsilon_n \sin \left(\frac{m\pi}{2} + (2n+1)w \right). \quad (12.5)$$

Formal versions. The displayed theorem is exactly `Fabius.prod_sin_two_pow_eq_thueMorse_sum` in `LacunaryProductToSum.lean`; its formal proof is precisely the specialization above. The same module records the two binary reflection laws

$$\text{wt}_2(2^m - 1 - n) + \text{wt}_2(n) = m, \quad \varepsilon_{2^m-1-n} = (-1)^m \varepsilon_n \quad (n < 2^m), \quad (12.6)$$

as `Fabius.binaryWeight_compl` and `Fabius.thueMorseSign_compl`, and the quarter-turn identity

$$(-1)^m \cos \left(\frac{m\pi}{2} - x \right) = \cos \left(\frac{m\pi}{2} + x \right) \quad (12.7)$$

for real x , as `Fabius.neg_one_pow_mul_cos_sub`.

Putting $w = z/2^m$ in (12.5) and reversing the finite product gives, for real z ,

$$\prod_{r=0}^m \sin \frac{z}{2^r} = \frac{1}{2^m} \sum_{n=0}^{2^m-1} \varepsilon_n \sin \left(\frac{\pi m}{2} + \frac{(2n+1)z}{2^m} \right). \quad (12.8)$$

The finite exponential identity also suggests the companion

$$\cos z \prod_{r=1}^m \sin \frac{z}{2^r} = \frac{1}{2^m} \sum_{n=0}^{2^m-1} \varepsilon_n \cos \left(\frac{\pi m}{2} + \frac{(2n+1)z}{2^m} \right). \quad (12.9)$$

where the product from $r = 1$ to 0 is empty when $m = 0$. Both formulas extend from real to complex z by the identity theorem.

Proof of the analytic companions. Set $\theta = 2z/2^m$ in (12.2), multiply by $i^m e^{iz/2^m}$, and simplify the phase. The result is

$$\sum_{n < 2^m} \varepsilon_n \exp\left(i \left(\frac{\pi m}{2} + \frac{(2n+1)z}{2^m} \right)\right) = 2^m e^{iz} \prod_{r=1}^m \sin \frac{z}{2^r}.$$

For real z , take imaginary and real parts. Both resulting identities are identities between entire functions of z , so the identity theorem extends them to every complex z . $\square$

Exact status boundary. Lean proves the raw finite exponential product for every complex parameter, together with its arbitrary affine translate. The complex half-angle and total sinc normalization are `Fabius.sum_thueMorseSign_cexp_eq_sin_prod` and `Fabius.thueMorseBlock_cexp_eq_q_sincPrefix` in `ThueMorseComplexProductBridge.lean`. Over $\mathbb{R}$, Lean proves (12.4), (12.5), and the cosine companion (12.9); the last is `Fabius.prod_sin_inv_two_pow_mul_cos_eq_thueMorse_sum` in `LacunaryProductToSumCosine.lean`. Thus the remaining packaging gap in this paragraph is only a pair of declarations stated directly in the two rescaled product-to-sum forms for arbitrary complex z . No raw complex transform, complex half-angle normalization, or real cosine companion remains unproved.

12.3 The exact dyadic discrete Fourier transform

Let

$$\text{DFT}_{2^m}[\varepsilon][k] = \sum_{n=0}^{2^m-1} \varepsilon_n \exp\left(-\frac{2\pi i k n}{2^m}\right), \quad 0 \leq k < 2^m. \quad (12.10)$$

Theorem 12.3 (Dyadic DFT). *For every k ,*

$$\text{DFT}_{2^m}[\varepsilon][k] = \prod_{j=0}^{m-1} \left(1 - \exp\left(-\frac{2\pi i k 2^j}{2^m}\right)\right). \quad (12.11)$$

If $m \geq 1$, this is zero for every even k . For odd k ,

$$\boxed{\text{DFT}_{2^m}[\varepsilon][k] = (2i)^m e^{-\pi i k (1-2^{-m})} \prod_{r=1}^m \sin \frac{\pi k}{2^r}.} \quad (12.12)$$

Thus

$$|\text{DFT}_{2^m}[\varepsilon][k]| = 2^m \prod_{r=1}^m \left| \sin \frac{\pi k}{2^r} \right| \quad (k \text{ odd}). \quad (12.13)$$

Proof. Evaluate (6.3) at $z = e^{-2\pi i k / 2^m}$. If k is even, the factor with $j = m-1$ is $1 - (-1)^k = 0$. For odd k , apply $1 - e^{-iu} = 2ie^{-iu/2} \sin(u/2)$ factorwise and sum the phases. $\square$

Formal versions. The product form (12.11) is the master product at $z = e^{-2\pi i k / 2^m}$; the closed sine form (12.12) is the specialization of `Fabius.sum_thueMorseSign_exp_eq_sin_prod` at $x = -2\pi k / 2^m$ (the atlas phase convention conjugates it). The vanishing at even frequencies is `Fabius.sum_thueMorseSign_mul_pow_eq_zero_of_pow_eq_one` (`ThueMorseRarefied.lean`). The inversion formula (12.14) is now formal as well: `Fabius.thueMorse_dft_inversion` (`ThueMorseFourierInversion.lean`) specializes a finite Fourier inversion theorem proved over an *arbitrary* field with a primitive root of unity (`Fabius.dft_inversion`, resting on the orthogonality `Fabius.sum_pow_eq_ite`) — a geometric-series fact, no integration.

For $m \geq 1$, Fourier inversion gives a root-of-unity formula for every individual term:

$$\varepsilon_n = \frac{1}{2^m} \sum_{\substack{0 \leq k < 2^m \\ k \text{ odd}}} \text{DFT}_{2^m}[\varepsilon][k] e^{2\pi i k n / 2^m}, \quad 0 \leq n < 2^m. \quad (12.14)$$

This is intentionally extravagant as an evaluator, but exact and natural for spectral calculations.

13 Walsh characters, Rademacher functions, Riesz products, and autocorrelation

13.1 A one-point Walsh spectrum

For m -bit vectors $a, b \in \mathbb{F}_2^m$, write $\langle a, b \rangle = \sum_j a_j b_j \pmod{2}$, and let $\mathbf{1} = (1, \dots, 1)$. For $n < 2^m$,

$$\varepsilon_n = (-1)^{\langle \mathbf{1}, \mathbf{b}(n) \rangle}. \quad (13.1)$$

Theorem 13.1 (Walsh transform is a point mass). *For every $a \in \mathbb{F}_2^m$,*

$$\sum_{n=0}^{2^m-1} \varepsilon_n (-1)^{\langle a, \mathbf{b}(n) \rangle} = \begin{cases} 2^m, & a = \mathbf{1}, \\ 0, & a \neq \mathbf{1}. \end{cases} \quad (13.2)$$

Proof. The summand is the character indexed by $a + \mathbf{1}$. A nontrivial character sums to zero over the finite Boolean group, while the trivial character sums to its order 2^m . $\square$

Thus the broad ordinary Fourier spectrum is a basis-change phenomenon: the same block is spectrally pure in the Walsh basis.

Formal versions. `ThueMorseWalsh.lean` proves (13.1) as `Fabius.thueMorseSign_eq_neg_one_pow_land` (the inner product realized as `wt2(a & n)`), the point mass (13.2) as `Fabius.sum_t hueMorseSign_mul_walsh`, and the underlying character orthogonality $\sum_{n < 2^m} \varepsilon_{a \& n} = 2^m \mathbf{1}_{a=0}$ as `Fabius.sum_neg_one_pow_binaryWeight_land`. The bit toolkit behind them — the all-ones digit pattern of $2^m - 1$ and the masking of two-power sums — is `Fabius.sum_range_two_pow`, `Fabius.testBit_two_pow_sub_one`, and `Fabius.land_sum_two_pow`.

13.2 Rademacher step functions

For $x \in [0, 1)$, define the right-continuous Rademacher functions

$$r_j(x) = (-1)^{\lfloor 2^{j+1}x \rfloor}. \quad (13.3)$$

If $0 \leq n < 2^m$ and $x \in [n/2^m, (n+1)/2^m)$, then the first m binary digits of x encode n , so

$$\varepsilon_n = \prod_{j=0}^{m-1} r_j(x). \quad (13.4)$$

The finite Thue–Morse block is therefore the sampling of one Walsh function on the dyadic partition.

13.3 Finite Riesz products

Define

$$\rho_m(x) = 2^{-m} |P_m(e^{2\pi i x})|^2. \quad (13.5)$$

Theorem 13.2 (Riesz-product density). *For every $m \geq 0$,*

$$\rho_m(x) = \prod_{j=0}^{m-1} (1 - \cos(2\pi 2^j x)). \quad (13.6)$$

Moreover,

$$\int_0^1 \rho_m(x) dx = 1, \quad (13.7)$$

so $\rho_m(x) dx$ is a probability measure.

Proof. Use $|1 - e^{i\theta}|^2 = 2(1 - \cos \theta)$ in (6.3). Parseval gives

$$\int_0^1 \rho_m(x) dx = 2^{-m} \sum_{n < 2^m} |\varepsilon_n|^2 = 1.$$

□

Formal versions. The factorization (13.6) is `Fabius.normSq_sum_thueMorseSign_exp` (`ThueMorseFourier.lean`) after normalizing by 2^m and rescaling the variable, with nonnegativity `Fabius.prod_two_sub_two_cos_nonneg`. The Parseval mass (13.7) is formal as `Fabius.integral_thueMorse_double_sum` (`ThueMorseParseval.lean`): the polynomial times its reflected conjugate integrates over a period to 2^m , by the character orthonormality `Fabius.integral_exp_two_pi_int`.

A second, entirely independent route to (13.7) – and to every Fourier coefficient of a lacunary Riesz product below its spectral floor – is formal in `LacunaryRieszIntegral.lean`, where nothing mentions Thue–Morse. For integer frequencies $m_0, m_1, \dots$, arbitrary real amplitudes a_j and phases φ_j , and an integer probe frequency K ,

$$\int_0^1 \cos(2\pi Kt + \psi) \prod_{j < n} (1 + a_j \cos(2\pi m_j t + \varphi_j)) dt = \begin{cases} \cos \psi, & K = 0, \\ 0, & K \neq 0, \end{cases} \quad (13.8)$$

under the hypothesis – exactly the one the peeling induction consumes – that

$$|K| + \sum_{i < j} m_i < m_j \quad \text{for every } j \quad (13.9)$$

(`Fabius.integral_cos_mul_rieszProduct_of_sum_lt`), so that no signed combination $\sum_j \eta_j m_j$ with $\eta_j \in \{-1, 0, 1\}$ can reach $-K$. Two corollaries carry the classical shapes: the K -free hypothesis $m_0 + \sum_{i < j} m_i \leq m_j$ with $|K| < m_0$ (`Fabius.integral_cos_mul_rieszProduct_of_add_sum_le`), and the Hadamard gap $2m_j \leq m_{j+1}$ with $|K| < m_0$ (`Fabius.integral_cos_mul_rieszProduct`, obtained from the previous one through `Fabius.add_sum_le_of_two_mul_le`). At $K = 0$ no headroom below the floor is needed, and plain super-increasingness $\sum_{i < j} m_i < m_j$ already gives the mean value 1 (`Fabius.integral_rieszProduct_of_superIncreasing`; the Hadamard-gap restatement, which still assumes $2m_j \leq m_{j+1}$ and $m_0 > 0$, is `Fabius.integral_rieszProduct`). Consequently (13.7) is the instance $m_j = 2^j$, $a_j = -1$, $\varphi_j = 0$. The companion exact identity $\int_0^1 \prod_{j < n} \sin^2(\pi m_j t) dt = 2^{-n}$ (`Fabius.integral_prod_sin_sq_of_superIncreasing`, with the geometric and dyadic specializations

`Fabius.integral_prod_sin_sq_pow` and `Fabius.integral_prod_sin_sq_two_pow`) is what the Fourier-decay audit of the Rvachev up-function uses. The offset m_0 in the super-increasing hypothesis cannot simply be deleted when $K \neq 0$ – the frequencies $3, 4, 8, 16, \dots$ with $K = 1$ resonate through $m_1 - m_0 = 1$ – and that counterexample, formerly recorded in prose in the module header, is now formalized in `RieszSharpness.lean`. The integral is evaluated exactly, $\int_0^1 \cos(2\pi t) (1 + \cos 6\pi t)(1 + \cos 8\pi t) dt = \frac{1}{4}$ (`Fabius.integral_cos_mul_sharpProduct`), and packaged as an existential refutation (`Fabius.exists_superIncreasing_probe_integral_ne`) in which two factors already suffice. What this establishes is precise and no more: plain super-increasingness together with $|K| < m_0$ does not suffice once $K \neq 0$. It is a single witness, and does not show the headroom hypothesis is the weakest K -free hypothesis yielding the conclusion.

The finite measures converge weakly to the Thue–Morse spectral measure. The formal notation

$$d\mu_{\text{TM}}(x) = \prod_{j \geq 0} (1 - \cos(2\pi 2^j x)) dx \quad (13.10)$$

means this weak limit, not pointwise multiplication of ordinary densities. The limiting measure is purely singular continuous, a deeper theorem of substitution spectral theory [14, 15].

13.4 Finite and limiting autocorrelation

For $m, k \geq 0$, define the dyadic average

$$\eta_m(k) = \frac{1}{2^m} \sum_{n=0}^{2^m-1} \varepsilon_n \varepsilon_{n+k}, \quad (13.11)$$

where the second index is not reduced modulo 2^m .

Proposition 13.3 (Finite autocorrelation recursion). *For $r \geq 0$,*

$$\eta_{m+1}(2r) = \eta_m(r), \quad (13.12)$$

$$\eta_{m+1}(2r+1) = -\frac{\eta_m(r) + \eta_m(r+1)}{2}. \quad (13.13)$$

Proof. Split the defining sum into even and odd n . For an even shift, both halves reduce to $\varepsilon_j \varepsilon_{j+r}$. For an odd shift, they reduce to $-\varepsilon_j \varepsilon_{j+r}$ and $-\varepsilon_j \varepsilon_{j+r+1}$, respectively. $\square$

Formal versions. `ThueMorseAutocorrelation.lean` proves the recursion (13.12)–(13.13) in unnormalized integer form ($A_m(k) = 2^m \eta_m(k)$): `Fabius.thueMorseAutocorrelation_succ_even` and `Fabius.thueMorseAutocorrelation_succ_odd`, with $A_m(0) = 2^m$ (`Fabius.thueMorseAutocorrelation_zero_shift`) and the closed shift values $3A_m(1) = -2^m - 2(-1)^m$ and $3A_{m+1}(2) = -2^{m+1} - 4(-1)^m$ (`Fabius.three_mul_thueMorseAutocorrelation_one`, `Fabius.three_mul_thueMorseAutocorrelation_two`) – the exact finite forms of $\eta(1) = \eta(2) = -\frac{1}{3}$.

The dyadic scaffolding of that recursion is removable. Nothing in the two-scale split uses that the window is a power of two, only that it is *even*, so the module proves the recursion first for the window sums $R_N(k) = \sum_{n < N} \varepsilon_n \varepsilon_{n+k}$ (`Fabius.thueMorseWindowAutocorrelation`) in the form

$$R_{2K}(2r) = 2R_K(r), \quad R_{2K}(2r+1) = -(R_K(r) + R_K(r+1)) \quad (13.14)$$

(`Fabius.thueMorseWindowAutocorrelation_two_mul_even`, `Fabius.thueMorseWindowAutocorrelation_two_mul_odd`), with $R_N(0) = N$ (`Fabius.thueMorseWindowAutocorrelation_zero_shift`) and the trivial bound $|R_N(k)| \leq N$ (`Fabius.abs_thueMorseWindowAutocorrelation_le`); the dyadic layer above is the case $N = 2^m$, read off through

the definitional identity `Fabius.thueMorseAutocorrelation_eq_window`, and its bound $|A_m(k)| \leq 2^m$ (`Fabius.abs_thueMorseAutocorrelation_le`) puts every normalized value in $[-1, 1]$. The normalized limits at the three shifts with closed forms are formal as well — $\eta_m(0) \rightarrow 1$, $\eta_m(1) \rightarrow -\frac{1}{3}$, $\eta_m(2) \rightarrow -\frac{1}{3}$ (`Fabius.tendsto_thueMorseAutocorrelation_zero_shift`, `Fabius.tendsto_thueMorseAutocorrelation_one`, `Fabius.tendsto_thueMorseAutocorrelation_two`); the existence of the limit at a *general* shift k is [theorem 13.4](#) below, now formal as well, while its spectral consequences remain analytic frontiers.

Theorem 13.4 (Limiting autocorrelation). *For every fixed $k \geq 0$, the limit*

$$\eta(k) = \lim_{m \rightarrow \infty} \eta_m(k) \quad (13.15)$$

exists and is characterized by

$$\eta(0) = 1, \quad (13.16)$$

$$\eta(2r) = \eta(r), \quad (13.17)$$

$$\eta(2r+1) = -\frac{\eta(r) + \eta(r+1)}{2}. \quad (13.18)$$

In particular,

$$\eta(1) = -\frac{1}{3}, \quad \eta(2) = -\frac{1}{3}, \quad \eta(3) = \frac{1}{3}, \quad \eta(5) = 0. \quad (13.19)$$

Proof. Induct on k . The case $k = 0$ is constant. For $k = 1$, (13.13) gives the contraction

$$\eta_{m+1}(1) = -\frac{1 + \eta_m(1)}{2},$$

whose fixed point is $-1/3$. If $k = 2r > 1$, convergence reduces to the smaller index r . If $k = 2r+1 > 1$, both r and $r+1$ are smaller than k , so the odd recurrence and the induction hypothesis give convergence. Pass to the limit and evaluate recursively. $\square$

Formal version. [Theorem 13.4](#) is formal in `ThueMorseAutocorrelationLimit.lean`, with exactly the proof above. The normalized sequence $\eta_m(k)$ is `Fabius.normalizedAutocorrelation`, and the two finite recursions hold *exactly* at every scale, $\eta_{m+1}(2r) = \eta_m(r)$ and $\eta_{m+1}(2r+1) = -(\eta_m(r) + \eta_m(r+1))/2$ (`Fabius.normalizedAutocorrelation_succ_even`, `Fabius.normalizedAutocorrelation_succ_odd`); the limit $\eta(k)$ is *defined* by the recursion (13.16)–(13.18) (`Fabius.limitingAutocorrelation`, with `Fabius.limitingAutocorrelation_two_mul` and `Fabius.limitingAutocorrelation_two_mul_add_one` as its laws), so the characterization is by construction and the theorem is the convergence (13.15), `Fabius.tendsto_normalizedAutocorrelation`, by strong induction on k with the two formal base cases $k = 0, 1$ of `ThueMorseAutocorrelation.lean`. The samples (13.19) are `Fabius.limitingAutocorrelation_one`, `Fabius.limitingAutocorrelation_two`, `Fabius.limitingAutocorrelation_three`, `Fabius.limitingAutocorrelation_five`.

The numbers $\eta(k)$ are the Fourier coefficients of μ_{TM} . Thus the substitution recursion, autocorrelation, and Riesz product are exact combinatorial, coefficient, and measure-theoretic forms of the same spectral object. Extending by evenness, $\eta(-k) = \eta(k)$, the first values are

$$1, -\frac{1}{3}, -\frac{1}{3}, \frac{1}{3}, -\frac{1}{3}, 0, \frac{1}{3}, 0, -\frac{1}{3}, \frac{1}{6}, 0, -\frac{1}{6}, \frac{1}{3}, -\frac{1}{6}, 0, \frac{1}{6}. \quad (13.20)$$

13.5 A Mahler equation for the autocorrelation series

The autocorrelation sequence has its own generating function, and the recursion of [theorem 13.4](#) compresses into a single functional equation of Mahler type.

Theorem 13.5 (Autocorrelation Mahler equation). *Let $C(z) = \sum_{n \geq 0} \eta(n)z^n$ for $|z| < 1$. Then*

$$\boxed{2zC(z) = 1 - (1 - z)^2C(z^2)}. \quad (13.21)$$

Proof. Split the series by parity of the index and insert (13.17) and (13.18):

$$\begin{aligned} C(z) &= \sum_{n \geq 0} \eta(2n)z^{2n} + \sum_{n \geq 0} \eta(2n+1)z^{2n+1} \\ &= C(z^2) - \frac{z}{2} \sum_{n \geq 0} \eta(n)z^{2n} - \frac{z}{2} \sum_{n \geq 0} \eta(n+1)z^{2n} \\ &= C(z^2) - \frac{z}{2}C(z^2) - \frac{1}{2z}(C(z^2) - 1). \end{aligned}$$

Multiply by $2z$ and collect the $C(z^2)$ terms into $-(1 - z)^2C(z^2)$. $\square$

Equation (13.21) places the autocorrelation beside E and T in the same Mahler framework; the square $(1 - z)^2$ reflects the fact that η averages products of two signs. Since the $\eta(k)$ are also the Fourier coefficients of the spectral measure, the equation is a coefficient-level renormalization law for that singular continuous measure: the measure has no closed-form distribution function, yet its Fourier data satisfies a two-term finite recursion.

Formal versions. `ThueMorseMahler.lean` proves a *finite-level* Mahler equation sharper than (13.21): with $C_m(z) = \sum_k A_m(k)z^k$ the unnormalized dyadic autocorrelation series of [proposition 13.3](#), $zC_{m+1}(z) = 2^m - (1 - z)^2C_m(z^2)$ holds exactly at every level (`Fabius.autocorrelationSeries_mahler`); dividing by 2^{m+1} recovers (13.21) in the limit.

14 Combinatorics and dynamics of the infinite word

The sequence read as an infinite binary word $\mathbf{t} = \tau(0)\tau(1)\tau(2) \cdots$ has a second life in combinatorics on words. A *factor* of $\mathbf{t}$ is a finite block $\tau(i)\tau(i+1) \cdots \tau(i+\ell-1)$ occurring at some position i ; the word-combinatorial theory studies which factors occur, how often, and with which repetitions. This section proves the elementary statements in full and states the two classical deep ones with references.

14.1 Uniform recurrence and letter frequencies

Recall from (2.4) that $\mathbf{t}$ is the fixed point, starting with 0, of the substitution $\mu : 0 \mapsto 01, 1 \mapsto 10$. Iterating, $\mathbf{t} = \lim_m \mu^m(0)$, and $\mu^m(0)$, $\mu^m(1)$ are the length- 2^m blocks $W_m, \overline{W}_m$ of (2.3), called the level- m *supertiles*.

Proposition 14.1 (Uniform recurrence). *The word $\mathbf{t}$ is uniformly recurrent: for every factor w there is a length $L(w)$ such that every window of $L(w)$ consecutive letters of $\mathbf{t}$ contains an occurrence of w .*

Proof. A factor w of length ℓ occurs inside some supertile $\mu^k(a)$ with $2^k \geq \ell$ and $a \in \{0, 1\}$. Both letters occur in $\mu(0) = 01$, so every supertile of level $k+1$ contains both $\mu^k(0)$ and $\mu^k(1)$; in particular every level- $(k+1)$ supertile contains w . Since $\mathbf{t}$ is tiled by consecutive level- $(k+1)$ supertiles, every window of length $L(w) = 2^{k+2}$ contains a complete one. $\square$

Uniform recurrence says every pattern that occurs once occurs with bounded gaps. In dynamical terms, the closure of the shift orbit of $\mathbf{t}$ is a minimal, uniquely ergodic system; that standard consequence of primitive constant-length substitutions [14] is what guarantees the existence of the limiting frequencies and autocorrelations used above. The two letters have frequency exactly $\frac{1}{2}$: at dyadic lengths this is the exact balance $A(2q) = q$ of (11.5), and general lengths deviate by at most one by [corollary 11.2](#).

Formal versions. `ThueMorseRecurrence.lean` proves [proposition 14.1](#) in a form sharper than stated: the recurrence bound is explicit and *linear* — a factor of length $\ell \leq 2^k$ occurs in every window of length $15 \cdot 2^k$ (`Fabius.thueMorseBit_uniformly_recurrent`), so every factor of length ℓ recurs with gap at most 30ℓ . The route is index arithmetic rather than supertiles: the full block-concatenation laws $\text{wt}_2(q \cdot 2^k + r) = \text{wt}_2(q) + \text{wt}_2(r)$ and $\tau(q \cdot 2^k + r) = \tau(q) \oplus \tau(r)$ (`Fabius.binaryWeight_block_concat`, `Fabius.thueMorseBit_block_concat`) reduce any window to its offset and two block bits, and all four two-letter factors occur in every window of fourteen letters (`Fabius.exists_thueMorseBit_pair`).

14.2 Aperiodicity

Proposition 14.2 (No eventual period). *The sequence (ε_n) is not eventually periodic. In particular it is not periodic.*

Proof. Suppose $p \geq 1$ were a period valid from some index N_0 on. For any $r \geq 0$, choose m with $2^m > r + p$ and $2^m + r \geq N_0$. Block concatenation (2.6) gives $\varepsilon_{2^m+r} = -\varepsilon_r$ and $\varepsilon_{2^m+r+p} = -\varepsilon_{r+p}$; eventual periodicity equates the left sides, so $\varepsilon_{r+p} = \varepsilon_r$ for every $r \geq 0$, and p is a global period. Now compare reflections: by (2.8), for every m with $2^m > p$,

$$-1 = \varepsilon_{2^m} = \varepsilon_{2^m-p} = \varepsilon_{(2^m-1)-(p-1)} = (-1)^m \varepsilon_{p-1}.$$

The right side changes sign between consecutive m , while the left side is constant — a contradiction. $\square$

Formal versions. [Proposition 14.2](#) is `Fabius.thueMorseSign_not_eventually_periodic` (`ThueMorseAperiodicity.lean`), but the formal proof no longer runs the reflection argument above: aperiodicity is now a *corollary of Thue’s theorem* ([theorem 14.3](#)), through a purely word-combinatorial implication that holds over an arbitrary alphabet and knows nothing about ε . If $f : \mathbb{N}_0 \rightarrow A$ is a sequence over any set A and satisfies $f(n+p) = f(n)$ for all $n \geq N$ with $p \geq 1$, then the window of length $2p+1$ starting at N is already an overlap of period p — every index $N+j$ with $j \leq p$ lies past the threshold, so $f(N+j) = f(N+p+j)$ (`Fabius.exists_overlap_of_eventually_periodic`); hence an overlap-free sequence over any alphabet admits no eventual period (`Fabius.not_eventually_periodic_of_overlap_free`, `ThueMorseOverlapFree.lean`). [Proposition 14.2](#) is that implication instantiated at `Fabius.thueMorseSign` and fed with `Fabius.thueMorseSign_overlap_free`; the zero-one counterpart `Fabius.thueMorseBit_not_eventually_periodic` is the same instantiation at `Fabius.thueMorseBit`. The dyadic block-flip identities on which the earlier self-contained proof ran are retained in `ThueMorseAperiodicity.lean` for their own sake — `Fabius.thueMorseSign_two_pow`, the block flip `Fabius.thueMorseSign_two_pow_add`, and the promotion of an eventual period to a global one (`Fabius.thueMorseSign_period_of_eventually`) — but they are no longer on the path to the theorem.

14.3 Overlap-freeness and the critical exponent

A *square* is a nonempty word repeated twice in immediate succession, ww ; a *cube* is www ; an *overlap* is a word of the form $axaxa$ with a a single letter and x a possibly empty word — equivalently, a square ww prolonged by one further letter of the same period. The *critical exponent* of an infinite word is the supremum, over its factors w , of $|w|/p$, where p is the smallest period of w and $|w|$ its length.

Theorem 14.3 (Thue). *The Thue–Morse word contains no overlap. Consequently it contains no cube, and no factor has exponent strictly greater than 2; since squares of arbitrarily great length do occur (apply μ^m to the square $\tau(1)\tau(2) = 11$), the critical exponent equals*

$$\boxed{2}, \tag{14.1}$$

and it is attained.

This is Thue’s foundational theorem [3]; a modern proof by desubstitution — a shortest overlap in $\mu(\mathbf{t})$ would force a shorter one in $\mathbf{t}$ after a finite boundary check — appears in [12]. Overlap-freeness is far stronger than the directly visible absence of the cubes 000 and 111, and it was the original reason for the word’s construction, decades before its automatic, spectral, and Diophantine roles emerged.

Formal versions. Thue’s theorem is no longer only a literature theorem here: `ThueMorseOverlapFree.lean` contains a complete formal proof, `Fabius.thueMorseBit_not_overlap` (with packaged forms `Fabius.thueMorseBit_overlap_free` and, for the signs, `Fabius.thueMorseSign_overlap_free`). The formal proof descends through the two-scale laws directly rather than by desubstitution: an even period halves by reading one parity class of the window (the seed fact is $\tau(2k) \neq \tau(2k+1)$, `Fabius.thueMorseBit_pair_ne`); an odd period forces $\tau(l) = \tau(l+1)$ along a run of consecutive l , which for $p \geq 5$ contradicts the absence of triples (`Fabius.thueMorseBit_no_triple`), while $p \in \{1, 3\}$ die by a Boolean analysis of the window. Cube-freeness is `Fabius.thueMorseBit_cube_free`, and the dyadic squares showing that the exponent 2 is attained are `Fabius.thueMorseBit_square_two_pow`. The module also isolates the step by which this theorem now *implies* [proposition 14.2](#): over an arbitrary alphabet an eventual period $p \geq 1$ from a threshold N on already exhibits an overlap of period p at N (`Fabius.exists_overlap_of_eventually_periodic`), so no overlap-free sequence over any alphabet is eventually periodic (`Fabius.not_eventually_periodic_of_overlap_free`).

14.4 Complement and reversal closure

Complementing every letter commutes with the substitution up to swapping the seeds $0 \leftrightarrow 1$; hence with every factor w the language of $\mathbf{t}$ also contains its bitwise complement $\bar{w}$. For reversal, the reflection formula (2.8) says that reading the length- 2^m prefix backwards reproduces it when m is even and complements it when m is odd — in signed terms this is the reciprocity of [proposition 6.4](#). Every factor sits inside some such prefix, so the factor language is closed under reversal as well. By unique ergodicity, complementary factors occur with equal frequencies; the exact prefix balance of [theorem 11.1](#) is the one-letter case.

Formal versions. Both closures are formal with fully explicit witnesses: `Fabius.thueMorseBit_exists_complement` places the complemented factor at a shift by a large power of two, and `Fabius.thueMorseBit_exists_reversal` places the reversed factor at the reflection of the window through an even-exponent dyadic block, where (2.8) carries the sign $+1$ (`ThueMorseOverlapFree.lean`).

14.5 Linear factor complexity

Let $p(\ell)$ denote the *factor complexity*: the number of distinct factors of length ℓ .

Proposition 14.4 (Elementary linear bound). *For every $\ell \geq 1$,*

$$\ell + 1 \leq p(\ell) < 8\ell. \quad (14.2)$$

Proof. Choose k minimal with $2^k \geq \ell$, so $2^k < 2\ell$. The word is a concatenation of level- k supertiles, so every length- ℓ factor lies inside a pair of consecutive supertiles $\mu^k(a)\mu^k(b)$. There are at most four choices of (a, b) and at most 2^k starting offsets, whence $p(\ell) \leq 4 \cdot 2^k < 8\ell$. The lower bound is the Morse–Hedlund theorem [19]: an infinite word with $p(\ell) \leq \ell$ for some ℓ is eventually periodic, which [proposition 14.2](#) forbids; hence $p(\ell) \geq \ell + 1$ for all ℓ . $\square$

The counting in that proof is wasteful: a window needs the second supertile only if it actually crosses the boundary between the two, and few offsets do. Bookkeeping the two cases separately sharpens the constant.

Proposition 14.5 (Sharpened linear bound). *For all k, ℓ with $\ell \leq 2^k$,*

$$p(\ell) \leq 2 \cdot 2^k + 2(\ell - 1), \quad (14.3)$$

and consequently, for every $\ell \geq 1$,

$$\boxed{p(\ell) \leq 6\ell - 4.} \quad (14.4)$$

Proof. Fix k with $\ell \leq 2^k$ and classify a starting position i by its offset $r = i \bmod 2^k$ in the aligned level- k block. If $r < 2^k - (\ell - 1)$ then $r + j < 2^k$ for every $j < \ell$, so the whole window lies inside that one block and is a function of the pair $(r, \tau(\lfloor i/2^k \rfloor))$: at most $2(2^k - (\ell - 1))$ windows arise this way. Otherwise r is one of the $\ell - 1$ offsets in $[2^k - \ell + 1, 2^k)$, the window straddles the boundary, and it is a function of the triple $(r, \tau(\lfloor i/2^k \rfloor), \tau(\lfloor i/2^k \rfloor + 1))$: at most $4(\ell - 1)$ windows. The two counts add to $2 \cdot 2^k + 2(\ell - 1)$, which is (14.3). Taking $k = \lceil \log_2 \ell \rceil$, so that $\ell \leq 2^k \leq 2\ell - 1$, turns it into $p(\ell) \leq 2(2\ell - 1) + 2(\ell - 1) = 6\ell - 4$. $\square$

For $\ell \geq 1$ the bound (14.4) implies (14.2), and (14.2) is retained only because it is the form the rest of the atlas quotes. Neither is the exact value: the exact piecewise-linear formula for $p(\ell)$ is a separate literature result, recalled but not reproved below.

Formal versions. Proposition 14.4 and proposition 14.5 are both fully formal in `ThueMorseComplexity.lean`, the sharp one being the primitive and the elementary one now its corollary. The complexity is `Fabius.thueMorseComplexity`, the cardinality of the set of length- ℓ windows. A window of length $\ell \leq 2^k$ is reconstructed from its offset in the aligned level- k block and the two adjacent block bits (`Fabius.thueMorseWindow_eq_reconstruct`), a statement that also records *where* the second bit is consulted — only for the indices j with $i \bmod 2^k + j \geq 2^k$ — which is what makes the case split of proposition 14.5 available. Splitting the image accordingly, *any* collection of windows contains at most $2 \cdot 2^k + 2(\ell - 1)$ distinct ones (`Fabius.card_image_thueMorseWindow_le'`), whence (14.3) (`Fabius.thueMorseComplexity_le'`) and, at $k = \lceil \log_2 \ell \rceil$, the boxed (14.4) (`Fabius.thueMorseComplexity_le_six_mul_sub_four`). The earlier chain survives with its signatures unchanged, each entry now a one-line consequence of its sharpened version: at most $4 \cdot 2^k$ windows (`Fabius.card_image_thueMorseWindow_le`), $p(\ell) \leq 4 \cdot 2^k$ (`Fabius.thueMorseComplexity_le`), and $p(\ell) < 8\ell$ (`Fabius.thueMorseComplexity_lt`), so nothing that quoted (14.2) needs revisiting. The lower bound `Fabius.le_thueMorseComplexity` is the Morse–Hedlund argument proved from scratch: complexity is monotone (`Fabius.thueMorseComplexity_mono`), and it grows *strictly* (`Fabius.thueMorseComplexity_lt_succ`) — were $p(\ell + 1) = p(\ell)$, restriction would be a cardinality-preserving surjection and hence injective, every length- ℓ factor would extend uniquely to the right, the finitely many windows would evolve deterministically, the pigeonhole principle would force two equal windows, and the word would be eventually periodic — contradicting the formal aperiodicity theorem. Uniform recurrence (`Fabius.thueMorseFactorSet_eq_image`) makes the factor set a computable finite image. Every bound holds in at least the stated form: the upper bounds `Fabius.thueMorseComplexity_le_six_mul_sub_four` and `Fabius.thueMorseComplexity_lt` need only $\ell \geq 1$, which is why the boxed lines (14.2) and (14.4) carry that hypothesis, while `Fabius.le_thueMorseComplexity` carries no side condition at all and holds at $\ell = 0$ too, where $p(0) = 1$ (`Fabius.thueMorseComplexity_zero`); the level form `Fabius.thueMorseComplexity_le'` needs only $\ell \leq 2^k$. No literature import is needed — and none of these statements asserts anything about the exact value of $p(\ell)$.

Linear complexity has two immediate consequences. The topological entropy, $\lim_{\ell} \ell^{-1} \log p(\ell)$, is zero. And the word is not *normal*: a normal binary word contains all 2^ℓ blocks of every length ℓ , which is incompatible with (14.2). Exact piecewise-linear formulas for $p(\ell)$ are known [12]; the elementary bound above suffices for every use in this atlas. The pointed contrast with section 15 below is worth savoring: the word itself is combinatorially rigid, yet its arithmetic subsequences behave pseudorandomly.

15 Rarefied sums, arithmetic progressions, and Newman’s phenomenon

Prefix sums of ε are bounded by [theorem 11.1](#), and sums over aligned dyadic blocks vanish. Restricting instead to an arithmetic progression destroys this balance in a striking way: along multiples of three one sign is *always* in excess, with a power-law surplus. The finite product P_m evaluated at roots of unity controls the whole phenomenon.

15.1 Root-of-unity filters at dyadic endpoints

For a modulus $q \geq 2$ and residue $0 \leq r < q$, define

$$A_{m,r}^{(q)} = \sum_{\substack{0 \leq n < 2^m \\ n \equiv r \pmod{q}}} \varepsilon_n. \quad (15.1)$$

Let $\zeta_q = e^{2\pi i/q}$, a primitive q -th root of unity. Averaging the geometric identity $\frac{1}{q} \sum_{\ell < q} \zeta_q^{(n-r)\ell} = \mathbf{1}_{q|n-r}$ against (6.3) gives the exact finite formula

$$A_{m,r}^{(q)} = \frac{1}{q} \sum_{\ell=0}^{q-1} \zeta_q^{-r\ell} P_m(\zeta_q^\ell) = \frac{1}{q} \sum_{\ell=0}^{q-1} \zeta_q^{-r\ell} \prod_{j=0}^{m-1} (1 - \zeta_q^{\ell 2^j}). \quad (15.2)$$

For $m \geq 1$ the term $\ell = 0$ vanishes, and the exponents $\ell 2^j$ move along the orbit of multiplication by two modulo q . For *even* moduli the phenomenon collapses: powers of two are eventually $0 \pmod{2^k}$, and for example $A_{m,0}^{(2)} = \sum_{k < 2^{m-1}} \varepsilon_{2k} = S(2^{m-1}) = 0$ for $m \geq 2$ by (11.2). Odd moduli are the interesting case: there the orbit of 2 is a fixed cycle and the products in (15.2) grow geometrically.

Formal versions. The identity (15.2) itself is now formal in `ThueMorseArithmeticFilter.lean`, in the multiplied form $q A_{m,r}^{(q)} = \sum_{\ell < q} \zeta_q^{-r\ell} \prod_{j < m} (1 - \zeta_q^{\ell 2^j})$ (`Fabius.thueMorse_rarefied_filter`). It specializes `Fabius.sum_filter_modEq_mul_natCast`, the filter proved for an *arbitrary* field carrying a primitive q -th root of unity and an *arbitrary* coefficient function — stated without division so that no invertibility of q is assumed; the progression indicator arises from the geometric orthogonality `Fabius.sum_pow_eq_ite` by exponent reduction and primitive-root injectivity. The even-modulus collapse $A_{m,0}^{(2)} = 0$ for $m \geq 2$ is `Fabius.sum_thueMorseSign_even_range`. The algebraic skeleton of the filter’s vanishing is `Fabius.sum_thueMorseSign_mul_pow_eq_zero_of_pow_eq_one` (`ThueMorseRarefied.lean`): over any commutative ring the signed block sum vanishes at any point with $z^{2^j} = 1$ for some $j < m$. Modulo three the filter is subsumed by the closed vectors below.

15.2 Complete evaluation modulo three

For $q = 3$ everything is explicit, because $2^j \equiv 1, 2, 1, 2, \dots \pmod{3}$ and, with $\omega = \zeta_3 = e^{2\pi i/3}$,

$$(1 - \omega)(1 - \omega^2) = 3. \quad (15.3)$$

Theorem 15.1 (Exact dyadic rarefaction modulo three). *For $u \geq 1$,*

$$(A_{2u,0}^{(3)}, A_{2u,1}^{(3)}, A_{2u,2}^{(3)}) = 3^{u-1} (2, -1, -1), \quad (15.4)$$

and for $u \geq 0$,

$$(A_{2u+1,0}^{(3)}, A_{2u+1,1}^{(3)}, A_{2u+1,2}^{(3)}) = 3^u (1, -1, 0). \quad (15.5)$$

Proof. At even level $m = 2u$, the factors of $P_m(\omega)$ pair into u copies of (15.3), so $P_{2u}(\omega) = P_{2u}(\omega^2) = 3^u$. At odd level one factor is left over: $P_{2u+1}(\omega) = 3^u(1 - \omega)$ and $P_{2u+1}(\omega^2) = 3^u(1 - \omega^2)$. Insert these values into (15.2) with $q = 3$ and simplify with $1 + \omega + \omega^2 = 0$. $\square$

Formal versions. `ThueMorseRarefied.lean` proves the closed vectors (15.4)–(15.5) as `Fabius.rarefiedSum_three_closed` – and does so without any root of unity: the one-step halving recursion $A_{m+1,r} = A_{m,2r \bmod 3} - A_{m,(2r+1) \bmod 3}$ (`Fabius.rarefiedSum_three_succ`; the hypothesis $r < 3$ is necessary, as the recursion is false for out-of-range residues) is iterated over paired levels by pure integer induction, so the filter argument is unnecessary for this particular modulo-three proof. The general root-of-unity filter itself is nevertheless formalized, in its division-free field form, by `Fabius.thueMorse_rarefied_filter` in `ThueMorseArithmeticFilter.lean`. Newman positivity at dyadic endpoints is `Fabius.rarefiedSum_three_zero_pos`, the full-block balance is `Fabius.sum_thueMorseSign_range_two_pow`, and the even-modulus collapse of the surrounding text is `Fabius.rarefiedSum_two_zero`.

The first vectors are

$$(1, -1, 0), \quad (2, -1, -1), \quad (3, -3, 0), \quad (6, -3, -3), \quad (9, -9, 0),$$

growing by a factor of three every two levels. Writing $X = 2^m$ for the endpoint, the excess on multiples of three is

$$A_{m,0}^{(3)} = c_m X^\alpha, \quad \alpha = \log_4 3 = \frac{\log 3}{\log 4} = 0.79248\dots, \quad (15.6)$$

where the constant c_m equals $\frac{2}{3}$ for even m and $\frac{1}{\sqrt{3}}$ for odd m : a clean two-scale oscillation around the power law.

15.3 Gelfond’s uniform exponent

The rarefied sums are finitely many samples of the trigonometric polynomial $p_m(x) = P_m(e^{2\pi i x})$, whose modulus is the sine product (12.3). How large can it get? Evaluating at $x = \frac{1}{3}$ gives, by (15.3),

$$\sup_{x \in \mathbb{R}} |p_m(x)| \geq |P_m(\omega)| = 3^{m/2} \quad (m \geq 0), \quad (15.7)$$

since the factors have modulus $|1 - \omega| = |1 - \omega^2| = \sqrt{3}$ in every position. (The attainment is formal: `Fabius.norm_sum_thueMorseSign_exp_two_pi_div_three` in `ThueMorseGelfond.lean` evaluates the modulus at the frequency $2\pi/3$ – the atlas’s $x = \frac{1}{3}$ – to exactly $(\sqrt{3})^m$ for every $m \geq 0$, not merely as a lower bound.) Gelfond’s theorem [28] – proved below – shows this is sharp up to a bounded factor:

$$\boxed{\sup_{x \in \mathbb{R}} |p_m(x)| \leq \frac{2}{\sqrt{3}} 3^{m/2} = \frac{2}{\sqrt{3}} (2^m)^{\log_4 3}}. \quad (15.8)$$

Proof of (15.8). The whole bound reduces to a single one-variable inequality. Write $\varphi(t) = 2|\sin t|$, so that $|p_m(x)| = \prod_{j < m} \varphi(2^j \pi x)$. For every real t ,

$$\varphi(t)^2 \varphi(2t) \leq 3\sqrt{3}, \quad (15.9)$$

because after squaring and substituting $u = \sin^2 t$ (so that $\sin^2 2t = 4u(1 - u)$) the claim reads $256 u^3(1 - u) \leq 27$, and

$$27 - 256u^3 + 256u^4 = (4u - 3)^2((4u + 1)^2 + 2) \geq 0,$$

with equality exactly at $u = \frac{3}{4}$ — the doubling-fixed Gelfond angle $t = \pi/3$, where $\varphi = \sqrt{3}$. Multiplying (15.9) along the doubling orbit, $\prod_{j < m-1} \varphi_j^2 \varphi_{j+1} \leq (3\sqrt{3})^{m-1}$ with $\varphi_j := \varphi(2^j t)$; the regrouping

$$\left(\prod_{j < m} \varphi_j \right)^3 = \varphi_0 \varphi_{m-1}^2 \prod_{j < m-1} (\varphi_j^2 \varphi_{j+1})$$

and the crude boundary bounds $\varphi_0 \leq 2$, $\varphi_{m-1}^2 \leq 4$ then give $(\prod_{j < m} \varphi_j)^3 \leq 8 \cdot 3^{3(m-1)/2} = \left(\frac{2}{\sqrt{3}} 3^{m/2}\right)^3$. Taking cube roots yields (15.8); the constant is Gelfond’s, and by (15.7) it cannot be lowered below 1. $\square$

Formal versions. The bound (15.8) is fully formal, with the proof above (ThueMorseGelfond.lean): the key inequality (15.9) is Fabius.two_abs_sin_sq_mul_le, proved by exactly the sum-of-squares certificate displayed; the telescoped product bound $\prod_{j < m} 2|\sin(2^j \theta)| \leq 2(\sqrt{3})^{m-1}$ is Fabius.gelfond_prod_le; and the exponential-sum form $|\sum_{n < 2^m} \varepsilon_n e^{inx}| \leq 2(\sqrt{3})^{m-1} = \frac{2}{\sqrt{3}} 3^{m/2}$, uniform in the real frequency x (the normalization of the frequency is immaterial to a uniform bound), is Fabius.gelfond_uniform_bound, riding the corpus sine-product identity Fabius.sum_thueMorseSign_exp_eq_sin_prod.

Decomposing an arbitrary range $[0, N)$ into at most one dyadic block per binary digit of N and applying (15.8) blockwise gives

$$\sup_{x \in \mathbb{R}} \left| \sum_{n=0}^{N-1} \varepsilon_n e^{2\pi i n x} \right| = \mathcal{O}(N^{\log_4 3}), \quad (15.10)$$

so $\log_4 3$ is the optimal uniform growth exponent of Thue–Morse exponential sums [28, 29]. At $x = 0$ the sum collapses to the bounded prefix values — the supremum and the special value differ by a full power of N .

Formal version. Equation (15.10) is formal with an explicit constant and for every N , not only asymptotically: Fabius.gelfond_general_bound in ThueMorseGelfond.lean gives $|\sum_{n < N} \varepsilon_n e^{inx}| \leq (\sqrt{3} + 1)(\sqrt{3})^{\lfloor \log_2 N \rfloor}$ uniformly in the real frequency x , by the blockwise decomposition described above (Fabius.norm_sum_thueMorseSign_exp_le_of_lt_two_pow), and Fabius.gelfond_general_bound_rpow rewrites the right side as $(\sqrt{3} + 1) N^{\log_4 3}$, which is the displayed $\mathcal{O}(N^{\log_4 3})$ with the implied constant made $\sqrt{3} + 1$.

15.4 A matrix recursion for odd moduli

For a fixed odd q , the vector $\mathbf{A}_m = (A_{m,r}^{(q)})_{0 \leq r < q}$ obeys a linear recursion. Splitting (15.1) by the parity of n , and writing 2^{-1} for the inverse of 2 modulo q (which exists because q is odd),

$$A_{m+1,r}^{(q)} = A_{m,2^{-1}r \bmod q}^{(q)} - A_{m,2^{-1}(r-1) \bmod q}^{(q)}, \quad (15.11)$$

since even $n = 2n' \equiv r$ forces $n' \equiv 2^{-1}r$ with $\varepsilon_n = \varepsilon_{n'}$, and odd $n = 2n' + 1 \equiv r$ forces $n' \equiv 2^{-1}(r-1)$ with $\varepsilon_n = -\varepsilon_{n'}$. Thus $\mathbf{A}_{m+1} = M_q \mathbf{A}_m$ for a fixed matrix M_q with one entry $+1$ and one entry -1 per row; diagonalizing M_q by the discrete Fourier transform on $\mathbb{Z}/q\mathbb{Z}$ reproduces (15.2). For $q = 3$ the eigenvalues have modulus $\sqrt{3}$, which is the growth rate seen in theorem 15.1.

Formal versions. The recursion (15.11) is Fabius.rarefiedSum_succ_of_odd (ThueMorseRarefied.lean), proved for every odd modulus with the halving realized by $2^{-1} = (q+1)/2$; the invertibility of doubling modulo an odd modulus is isolated as Fabius.two_mul_mod_iff. Modulo three it specializes to Fabius.rarefiedSum_three_succ.

15.5 The global Newman theorem and the arithmetic frontier

Theorem 15.1 evaluates the excess at dyadic endpoints. Newman’s theorem extends the positivity to *every* endpoint. The positivity assertion of Newman’s theorem is proved below by a short base-four induction — and is formalized; the quantitative refinements, and the deeper equidistribution results that follow, are literature theorems, stated here to locate the atlas’s exact identities inside modern analytic number theory.

Theorem 15.2 (Newman). *For every real $X \geq 1$,*

$$N_3(X) := \sum_{\substack{0 \leq n < X \\ 3|n}} \varepsilon_n > 0. \quad (15.12)$$

Moreover $\frac{1}{20}X^{\log_4 3} \leq N_3(X) \leq 5X^{\log_4 3}$ for $X \geq 2$.

Proof of the positivity. It suffices to treat integer cutoffs $N \geq 1$. Write $T_r(N) = \sum_{n < N, n \equiv r(3)} \varepsilon_n$ and $E(N) = \sum_{n < N} \varepsilon_n$; consecutive pairs cancel, so $E(2M) = 0$ and $E(2M+1) = \varepsilon_{2M} = \pm 1$ — in particular $|E| \leq 1$ everywhere. The parity split $\varepsilon_{2j} = \varepsilon_j, \varepsilon_{2j+1} = -\varepsilon_j$ gives the one-step system

$$T_0(2M) = T_0(M) - T_1(M), \quad T_1(2M) = T_2(M) - T_0(M),$$

and composing two steps and eliminating $T_1(M) + T_2(M) = E(M) - T_0(M)$ collapses to

$$T_0(4M+s) = 3T_0(M) - E(M) + c_s(M), \quad s \in \{0, 1, 2, 3\},$$

where the boundary correction $c_s(M)$ collects at most three single signs: ε_M from the term $n = 4M$ (alive only when $M \equiv 0$), $-\varepsilon_M$ from the odd term of the second halving (alive only when $M \equiv 2$), and $-\varepsilon_M$ from the term $n = 4M+2$ (alive only when $M \equiv 1$, since $4M+2 \equiv M+2$, and $\varepsilon_{4M+2} = -\varepsilon_M$). The three supports are pairwise disjoint residue classes of M , so $|c_s(M)| \leq 1$. By strong induction $T_0(M) \geq 1$ for all $M \geq 1$ — the cases $N \leq 4$ are read off the signs $+, -, -, +$ — whence $T_0(4M+s) \geq 3 \cdot 1 - 1 - 1 = 1$: the three-fold amplification absorbs all boundary noise. The quantitative bounds are Newman’s refinement of the same induction and are not reproved here. $\square$

Among multiples of three below any cutoff, integers of even binary weight always outnumber those of odd binary weight [30]. Coquet determined the precise structure: $N_3(X) = X^{\log_4 3} \Psi(\log_4 X) + \mathcal{O}(1)$ with Ψ a continuous, nowhere-differentiable periodic function [31].

Formal versions. The boxed positivity (15.12) is fully formal, with the proof above (Fabius.newman_positivity, ThueMorseNewman.lean): the residue-class sign sums are Fabius.thueMorseResidueSum (defined for every modulus), the parity system is Fabius.thueMorseResidueSum_three_zero_two_mul and its companion, the three sums recombine to the prefix sum by Fabius.thueMorseResidueSum_three_add, and the prefix-sum bound $|E| \leq 1$ is Fabius.abs_sum_thueMorseSign_le_one. The quantitative bounds $\frac{1}{20}X^{\log_4 3} \leq N_3(X) \leq 5X^{\log_4 3}$ and Coquet’s oscillation remain literature imports. Rarefactions along other progressions and for generalized digit characters were analyzed by Drmota and Skalba [32] and Drmota and Stoll [33]; a recent automata-theoretic treatment mechanizes such evaluations [34].

Beyond arithmetic progressions, the same digit-parity character is now understood along much thinner sequences. Mauduit and Rivat proved Gelfond’s conjecture that the parity is equidistributed along the primes: $\sum_{p \leq x} \varepsilon_p = o(\pi(x))$, the sum running over primes with $\pi(x)$ their count [35]. Drmota, Mauduit, and Rivat proved that along the squares the word $(\tau(n^2))_{n \geq 0}$ is *normal* — every binary block of length L occurs with limiting frequency 2^{-L} [36] — in sharp contrast with the non-normality of the word itself (proposition 14.4). Spiegelhofer established a level-of-distribution result covering $\tau(\lfloor n^c \rfloor)$ for $1 < c < 2$ [37] and the one-symbol balance along the cubes [38]. The deterministic word is combinatorially rigid, yet its polynomial and prime subsequences behave like fair coin flips: that tension is the modern face of Gelfond’s problems, and every one of these theorems begins from the finite product and carry structure catalogued in this atlas.

16 Boundary flatness, shifted Dirichlet series, and Mellin continuation

This section develops an analytic consequence not organized systematically in the four drafts. The lacunary product is extraordinarily flat as its argument approaches one from inside the unit disk. That flatness removes every Mellin singularity and gives an entire continuation of a natural family of Thue–Morse Dirichlet series.

16.1 Quadratic logarithmic decay near one

For $t > 0$, set

$$\mathcal{E}(t) = E(e^{-t}) = \prod_{j \geq 0} (1 - e^{-2^j t}). \quad (16.1)$$

Every factor lies in $(0, 1)$, so $\mathcal{E}(t) > 0$.

Theorem 16.1 (Boundary flatness). *As $t \downarrow 0$,*

$$\log \mathcal{E}(t) = -\frac{\log^2(1/t)}{2 \log 2} + \mathcal{O}(\log(1/t)). \quad (16.2)$$

Consequently, for every $A > 0$,

$$\mathcal{E}(t) = \mathcal{O}_A(t^A) \quad (t \downarrow 0). \quad (16.3)$$

Proof. Assume $0 < t < 1/2$, put $L = \log(1/t)$, and let $m = \lfloor \log_2(1/t) \rfloor$. Then $2^m t \in (1/2, 1]$. For $0 \leq j < m$, one has $0 < 2^j t \leq 1/2$, and the elementary bounds

$$\frac{x}{2} \leq 1 - e^{-x} \leq x \quad (0 < x \leq 1) \quad (16.4)$$

show that

$$\log(1 - e^{-2^j t}) = \log t + j \log 2 + \mathcal{O}(1), \quad (16.5)$$

with an absolute $\mathcal{O}(1)$. Summing gives

$$\sum_{j=0}^{m-1} \log(1 - e^{-2^j t}) = m \log t + \frac{m(m-1)}{2} \log 2 + \mathcal{O}(m). \quad (16.6)$$

For the tail, write $2^m t = u \in (1/2, 1]$. Then

$$\sum_{j=m}^{\infty} \left| \log(1 - e^{-2^j t}) \right| = \sum_{r=0}^{\infty} \left| \log(1 - e^{-2^r u}) \right| = \mathcal{O}(1) \quad (16.7)$$

uniformly in u , because the terms decay doubly exponentially. Finally $m = L/\log 2 + \mathcal{O}(1)$; substitution in (16.6) yields (16.2). The negative quadratic term dominates $-A \log(1/t)$ for every fixed A , proving (16.3). $\square$

Remark 16.2. Sharper asymptotic expansions contain a bounded log-periodic fluctuation in $\log_2(1/t)$. The elementary estimate above is deliberately limited to the part needed for Mellin continuation: super-polynomial flatness.

Formal versions. [Theorem 16.1](#) is formal in effective two-sided form (`ThueMorseBoundaryFlatness.lean`). The product $\mathcal{E}$ is `Fabius.lacunaryExpProduct`, identified with the exponential of its log series and hence positive and at most one. The key structural fact is the *self-similarity* $\mathcal{E}(t) = \prod_{j < m} (1 - e^{-2^j t}) \cdot \mathcal{E}(2^m t)$ (`Fabius.lacunaryExpProduct_eq_prod_mul`); combining it

with the reusable sandwich $x/2 \leq 1 - e^{-x} \leq x$ on $[0, 1]$ (`Fabius.half_le_one_sub_exp_neg`, `Fabius.one_sub_exp_neg_le` – the lower half proved from $e^{-x} \leq 1/(1+x)$, no series expansion) gives, for every level m ,

$$2^{\binom{m}{2}}(t/2)^m \mathcal{E}(\tfrac{1}{2}) \leq \mathcal{E}(t) \leq 2^{\binom{m}{2}}t^m,$$

the lower bound under the dyadic matching $\frac{1}{2} \leq 2^mt \leq 2$ (`Fabius.le_lacunaryExpProduct`, `Fabius.lacunaryExpProduct_le`). Taking logarithms at the matching level recovers (16.2) with explicit constants, and the upper bound alone is (16.3) in effective form. Monotonicity in t (`Fabius.lacunaryExpProduct_mono`) and summability of the factor logarithms (`Fabius.summable_log_one_sub_exp`) are proved along the way.

16.2 The shifted Dirichlet family

For $a > 0$, define initially in the half-plane $\Re s > 0$

$$D(s, a) = \sum_{n=0}^{\infty} \frac{\varepsilon_n}{(n+a)^s}. \quad (16.8)$$

The principal real logarithm is used in $(n+a)^{-s}$.

Proposition 16.3 (Conditional convergence). *The series (16.8) converges locally uniformly in $\Re s > 0$, and hence defines a holomorphic function there.*

Proof. The partial sums $S(N) = \sum_{n < N} \varepsilon_n$ are bounded by one. On a compact set $\Re s \geq \sigma > 0$, the first difference of $(n+a)^{-s}$ is uniformly $\mathcal{O}(n^{-\sigma-1})$. Summation by parts therefore gives uniform convergence of the tails. $\square$

Formal versions. The real regime of the family is formal in `ThueMorseDirichlet.lean`: for every $\sigma > 0$ and $a > 0$ the partial sums of $D(\sigma, a)$ form a Cauchy sequence by Dirichlet’s test – the antitone shifted powers tend to zero and the sign prefix sums are bounded by one – (`Fabius.dirichletPartial_cauchy`), and the limit function `Fabius.thueMorseDirichlet` is their limit (`Fabius.tendsto_thueMorseDirichlet`); this is [proposition 16.3](#) along the real axis. The dyadic equation below is `Fabius.thueMorseDirichlet_dyadic` on that same real ray – for every real $\sigma > 0$ and $a > 0$ – proved by interleaving the even-length partial sums through $\varepsilon(2n) = \varepsilon(n)$, $\varepsilon(2n+1) = -\varepsilon(n)$ and passing to limits along the doubled subsequence; the complex form is recorded after [theorem 16.4](#) below. The Mellin representation itself is now formal in the absolutely convergent real regime (`Fabius.thueMorseDirichlet_mellin`, `ThueMorseMellin.lean`): for $\sigma > 1$ and $a > 0$, $\Gamma(\sigma)D(\sigma, a) = \int_0^\infty t^{\sigma-1}e^{-at}\mathcal{E}(t) dt$ – the integrand is rewritten through the analytic infinite-product identity, the sum and integral are interchanged by absolute convergence (the norm integrals are exactly the Gamma values, summable as a shifted p -series), and each term is a Gamma integral. For $\sigma > 1$ the conditionally defined D is identified with the absolute sum (`Fabius.thueMorseDirichlet_eq_tsum`). The entire continuation itself is now formal (`ThueMorseEntireContinuation.lean`): the function $s \mapsto \Gamma(s)^{-1} \int_0^\infty t^{s-1}e^{-at}\mathcal{E}(t) dt$ is differentiable on all of $\mathbb{C}$ (`Fabius.dirichletMellinContinuation_differentiable`) – the effective flatness bound supplies an $\mathcal{O}(t^{-b})$ estimate at zero for every b (`Fabius.mellinKernel_isBigO_rpow`), precisely realizing the maxim that boundary flatness removes every Mellin singularity, while $1/\Gamma$ is entire – and it agrees with $D(\sigma, a)$ on the ray $\sigma > 1$ (`Fabius.dirichletMellinContinuation_eq`). Only the locally-uniform holomorphy of the conditionally convergent series in the strip $0 < \Re s \leq 1$ (superseded by the continuation) remains informal.

Splitting even and odd indices gives an exact parameter equation.

Theorem 16.4 (Dyadic Dirichlet equation). *For $a > 0$ and $\Re s > 0$,*

$$D(s, a) = 2^{-s} \left(D\left(s, \frac{a}{2}\right) - D\left(s, \frac{a+1}{2}\right) \right). \quad (16.9)$$

Proof. Use $\varepsilon_{2n} = \varepsilon_n$ and $\varepsilon_{2n+1} = -\varepsilon_n$:

$$\begin{aligned} D(s, a) &= \sum_{n \geq 0} \frac{\varepsilon_n}{(2n + a)^s} - \sum_{n \geq 0} \frac{\varepsilon_n}{(2n + 1 + a)^s} \\ &= 2^{-s} \left(D\left(s, \frac{a}{2}\right) - D\left(s, \frac{a+1}{2}\right) \right). \end{aligned}$$

□

Formal versions. [Theorem 16.4](#) is now formal at every complex s , not only along the real ray. The carrier is the entire continuation of [theorem 16.5](#), formalized as `Fabius.dirichletMellinContinuation`: for every $a > 0$ and every $s \in \mathbb{C}$,

$$\tilde{D}(s, a) = 2^{-s} \left(\tilde{D}\left(s, \frac{a}{2}\right) - \tilde{D}\left(s, \frac{a+1}{2}\right) \right) \quad (16.10)$$

(`Fabius.dirichletMellinContinuation_dyadic`, `ThueMorseEntireContinuation.lean`), where $\tilde{D}$ denotes that continuation. The proof is the identity theorem, and it is the reason the entirety layer was built: both sides of (16.10) are entire — the left by `Fabius.dirichletMellinContinuation_differentiable` at a , the right by the same at $a/2$ and $(a+1)/2$ multiplied by the entire $s \mapsto 2^{-s}$ — and they agree on the ray $(1, \infty) \subseteq \mathbb{R} \subseteq \mathbb{C}$ (`Fabius.dirichletMellinContinuation_dyadic_ofReal`, itself `Fabius.thueMorseDirichlet_dyadic` transported through `Fabius.dirichletMellinContinuation_eq`); that ray accumulates at $s = 2$, and $\mathbb{C}$ is connected. Note what is and is not delivered: (16.10) is unconditional in s , while its identification with (16.9) — that is, $\tilde{D} = D$ — is formal only where `Fabius.dirichletMellinContinuation_eq` applies, namely on the real ray $\sigma > 1$; off that ray the identification still rests on the informal half of [proposition 16.3](#).

16.3 Mellin representation and entire continuation

Theorem 16.5 (Mellin integral). *For $a > 0$, initially when $\Re s > 1$,*

$$D(s, a) = \frac{1}{\Gamma(s)} \int_0^\infty t^{s-1} e^{-at} E(e^{-t}) dt. \quad (16.11)$$

The integral on the right is an entire function of s . Therefore (16.11) continues $D(s, a)$ to an entire function of $s \in \mathbb{C}$.

Proof. For $\Re s > 1$, absolute convergence permits termwise use of

$$(n + a)^{-s} = \frac{1}{\Gamma(s)} \int_0^\infty t^{s-1} e^{-(n+a)t} dt.$$

Summing $\varepsilon_n e^{-nt}$ gives $E(e^{-t})$.

At infinity, $E(e^{-t}) = 1 + \mathcal{O}(e^{-t})$, and the factor e^{-at} gives exponential decay. At zero, [theorem 16.1](#) gives decay faster than every positive power of t . Thus the integral and all of its s -derivatives converge uniformly on compact subsets of the entire s -plane. Multiplication by the entire function $1/\Gamma(s)$ preserves entire-ness. Equality on $\Re s > 1$, together with [proposition 16.3](#), gives the continuation by the identity theorem. □

Corollary 16.6 (Zeros at nonpositive integers). *For every $a > 0$ and every integer $r \geq 0$,*

$$D(-r, a) = 0. \quad (16.12)$$

Proof. The Mellin integral in (16.11) is finite at $s = -r$, while $1/\Gamma(s)$ has a zero there. □

These zeros are the analytic continuation of the finite Prouhet phenomenon: negative integer values formally ask for signed polynomial moments, and the flat Mellin kernel makes all of them vanish after continuation.

Formal versions. [Corollary 16.6](#) is now formal (`ThueMorseEntireContinuation.lean`): `Fabius.dirichletMellinContinuation_neg_natCast` gives $D(-r, a) = 0$ for every integer $r \geq 0$, with the origin recorded separately as `Fabius.dirichletMellinContinuation_zero`; neither statement needs any hypothesis on a , since the vanishing is supplied by the zero of $1/\Gamma$ at $s = -r$. What makes that argument legitimate is the stronger, *un-damped* statement – and it too is formal – that the Mellin factor alone is entire, with no Γ -factor damping it (`Fabius.mellin_mellinKernel_differentiable`, for $a > 0$): boundary flatness leaves no pole at $s = -r$ that could cancel that zero.

16.4 The harmonic kernel

At $s = 1$, the Mellin integral can be moved to the unit interval.

Corollary 16.7 (Dirichlet integral). *For $a > 0$,*

$$H(a) := D(1, a) = \sum_{n \geq 0} \frac{\varepsilon_n}{n + a} \quad (16.13)$$

$$= \int_0^1 x^{a-1} E(x) dx = \int_0^1 x^{a-1} \prod_{j \geq 0} (1 - x^{2^j}) dx > 0. \quad (16.14)$$

It satisfies

$$H(a) = \frac{1}{2} H\left(\frac{a}{2}\right) - \frac{1}{2} H\left(\frac{a+1}{2}\right). \quad (16.15)$$

Proof. Set $s = 1$ in (16.11) and substitute $x = e^{-t}$. The integrand is positive on $(0, 1)$. Equation (16.15) is (16.9) at $s = 1$. $\square$

Formal status. The series (16.13) is the value at $\sigma = 1$ of `Fabius.thueMorseDirichlet` in `ThueMorseDirichlet.lean`, the limit of the partial sums `Fabius.dirichletPartial`, whose conditional convergence for every $\sigma > 0$, $a > 0$ is `Fabius.tendsto_thueMorseDirichlet`; the dyadic law (16.15) is the case $\sigma = 1$ of `Fabius.thueMorseDirichlet_dyadic`, where the factor $2^{-\sigma}$ becomes the $\frac{1}{2}$ displayed. The integral representation (16.14) and the positivity $H(a) > 0$ are not formalized: they rest on the Mellin identity (16.11), whose sum–integral interchange is conditional rather than dominated on $(0, 1)$.

17 Infinite products with Thue–Morse exponents

17.1 Pair reduction

Suppose $R(n) \neq 0$ and the products converge. The even–odd recurrence gives the general reduction

$$\prod_{n \geq 0} R(n)^{\varepsilon_n} = \prod_{q \geq 0} \left(\frac{R(2q)}{R(2q+1)} \right)^{\varepsilon_q}. \quad (17.1)$$

A factor merely equal to $1 + \mathcal{O}(1/n)$ is often converted into a paired factor $1 + \mathcal{O}(1/n^2)$, explaining the good convergence of balanced rational products.

17.2 A two-parameter master product

For $a, b > 0$, define

$$G(a, b) = \prod_{n=0}^{\infty} \left(\frac{n+a}{n+b} \right)^{\varepsilon_n}. \quad (17.2)$$

The logarithm converges by summation by parts: the Thue–Morse partial sums are bounded and the first difference of $\log((n+a)/(n+b))$ is $\mathcal{O}(n^{-2})$.

Theorem 17.1 (Cocycle and dyadic functional equation). *For positive parameters,*

$$G(a, a) = 1, \quad G(a, b)G(b, c) = G(a, c), \quad G(b, a) = G(a, b)^{-1}, \quad (17.3)$$

$$G(a, b) = \frac{G(a/2, b/2)}{G((a+1)/2, (b+1)/2)}. \quad (17.4)$$

Proof. The cocycle identities are termwise. Pair indices $2n$ and $2n+1$:

$$\begin{aligned} & \left(\frac{2n+a}{2n+b} \right)^{\varepsilon_n} \left(\frac{2n+1+a}{2n+1+b} \right)^{-\varepsilon_n} \\ &= \left[\frac{n+a/2}{n+b/2} \frac{n+(b+1)/2}{n+(a+1)/2} \right]^{\varepsilon_n}. \end{aligned}$$

Multiply over n . □

Formal versions. [Theorem 17.1](#) is formal in logarithmic form (`ThueMorseMasterProduct.lean`): the sign-weighted log-series $L(a, b) = \sum \varepsilon_n \log((n+a)/(n+b))$ converges for all positive parameters by Dirichlet’s test (`Fabius.mpLog_cauchy`); the cocycle identities are `Fabius.mpLimit_self`, `Fabius.mpLimit_symm`, and `Fabius.mpLimit_cocycle`; the boxed dyadic functional equation is `Fabius.mpLimit_dyadic`, proved by the exact even-length interleaving rather than by rearrangement; and the product itself converges to $\exp L(a, b)$ (`Fabius.tendsto_masterProduct`).

The Dirichlet continuation identifies this product as a zeta-regularized derivative difference without any extra regularization being required.

Corollary 17.2 (Dirichlet derivative bridge). *For $a, b > 0$,*

$$\log G(a, b) = D'(0, b) - D'(0, a). \quad (17.5)$$

Proof. The difference $D(s, a) - D(s, b)$ can be differentiated termwise at $s = 0$, because the resulting log-ratio series converges by summation by parts. Its derivative is

$$\sum_{n \geq 0} \varepsilon_n \log \frac{n+b}{n+a} = -\log G(a, b).$$

Rearrange. □

Remark 17.3 (A Fubini-free proof of the bridge). The sketch above can be completed without any uniform summation-by-parts estimate; the following route is the one the formalization takes. First, each term of the master log-series is a Frullani integral,

$$\varepsilon_n \log \frac{n+a}{n+b} = \varepsilon_n \int_0^\infty \frac{e^{-(n+b)t} - e^{-(n+a)t}}{t} dt,$$

and summing over $n < 2^m$ factors the signed exponentials into the *finite* lacunary product:

$$\sum_{n < 2^m} \varepsilon_n \log \frac{n+a}{n+b} = \int_0^\infty \frac{e^{-bt} - e^{-at}}{t} \prod_{j < m} (1 - e^{-2^j t}) dt.$$

The partial products lie in $[0, 1]$, so dominated convergence — with the integrable dominator $|a - b|e^{-\min(a,b)t}$, a consequence of $1 - e^{-x} \leq x$ — passes to the limit along the dyadic subsequence, along which the conditionally convergent series already converges:

$$\log G(a, b) = \int_0^\infty \frac{e^{-bt} - e^{-at}}{t} E(e^{-t}) dt. \quad (17.6)$$

Second, the factorization $\Gamma(s)^{-1} = s\Gamma(s+1)^{-1}$ — an identity of entire functions, valid at the poles as well — gives $(1/\Gamma)'(0) = 1$ and $\Gamma(0)^{-1} = 0$, so by the product rule the entire continuation $s \mapsto \Gamma(s)^{-1} \int_0^\infty t^{s-1} e^{-at} E(e^{-t}) dt$ of $D(s, a)$ constructed above has derivative at $s = 0$ equal to the Mellin value $\int_0^\infty t^{-1} e^{-at} E(e^{-t}) dt$ — convergent, since boundary flatness gives $E(e^{-t}) \leq t$. Subtracting the two parameters recovers (17.6).

Formal versions. [Corollary 17.2](#) is fully formal along the route of [Remark 17.3](#) (`ThueMorseGDirichlet.lean`). The Frullani integral $\int_0^\infty (e^{-pt} - e^{-qt})t^{-1} dt = \log(q/p)$ is proved from scratch by differentiation under the integral sign — the $1/t$ cancels exactly against the derivative of the exponential — followed by the fundamental theorem of calculus in the parameter (`Fabius.frullani_exp`, `FrullaniIntegral.lean`; the lemma is not in `Mathlib`). The closed integral form (17.6) is `Fabius.mpLimit_eq_integral`; the derivative $(1/\Gamma)'(0) = 1$ is `Fabius.hasDerivAt_Gamma_inv_zero`, riding `Mathlib`’s entire-function factorization $\Gamma(s)^{-1} = s\Gamma(s+1)^{-1}$; the derivative of the entire continuation at zero is `Fabius.hasDerivAt_dirichletMellinContinuation_zero`; and the boxed bridge (17.5) is `Fabius.mpLimit_eq_deriv_sub`, with $D(\cdot, c)$ realized by its entire continuation.

17.3 The Woods–Robbins product

Theorem 17.4 (Woods–Robbins).

$$\prod_{n=0}^{\infty} \left(\frac{2n+1}{2n+2} \right)^{\varepsilon_n} = \frac{1}{\sqrt{2}}. \quad (17.7)$$

Proof. Let

$$A = \prod_{n \geq 0} \left(\frac{2n+1}{2n+2} \right)^{\varepsilon_n}, \quad B = \prod_{n \geq 1} \left(\frac{2n}{2n+1} \right)^{\varepsilon_n}. \quad (17.8)$$

Then

$$AB = \frac{1}{2} \prod_{n \geq 1} \left(\frac{n}{n+1} \right)^{\varepsilon_n}. \quad (17.9)$$

Split the product on the right into odd and even n . The odd part is A^{-1} because $\varepsilon_{2j+1} = -\varepsilon_j$, and the even part is B because $\varepsilon_{2j} = \varepsilon_j$. Hence

$$AB = \frac{1}{2} A^{-1} B. \quad (17.10)$$

Cancel B , obtaining $A^2 = 1/2$; positivity gives the stated root. $\square$

Formal versions. [Theorem 17.4](#) is fully formal (`Fabius.woods_robbins`, `ThueMorseWoodsRobbins.lean`), with the pairing argument carried out rigorously in logarithmic form: the three sign-weighted log-series converge by Dirichlet’s test (a reusable guarded variant handles weights antitone only from 1 on); the termwise combination $L_A + L_B = -\log 2 + L_C$ and the exact even-length interleaving $L_C(2N) = L_B(N) - L_A(N)$ — the formal counterpart of the odd/even split, justified along the doubled subsequence rather than by rearrangement — force $L_A = -\frac{1}{2} \log 2$ (`Fabius.tendsto_wra`), and exponentiation evaluates the partial products. That final exponentiation

is itself a reusable module, `RealZPowProduct.lean`, in which nothing mentions Thue–Morse: a finite product of integer powers of positive reals is the exponential of the signed logarithm sum (`Fabius.prod_zpow_eq_exp_sum`), so a convergent logarithm sum exponentiates to a convergent product along an arbitrary filter (`Fabius.tendsto_prod_zpow_of_tendsto_sum`), and the closed values come from `Fabius.exp_neg_mul_log` with its two specializations `Fabius.exp_neg_log_div_two` ($e^{-\frac{1}{2} \log b} = 1/\sqrt{b}$) and `Fabius.exp_neg_three_halves_log`. Every product evaluation below reuses it. No convergence step is left informal.

A second classical evaluation is

$$\prod_{n=0}^{\infty} \left(\frac{4n+1}{4n+3} \right)^{\varepsilon_n} = \frac{1}{2}. \quad (17.11)$$

It is the specialization $G(1/4, 3/4)$; it appears in the broader Thue–Morse gamma-product calculus of Allouche, Riasat, and Shallit [16], but it also follows from the Woods–Robbins evaluation alone, by one more parity split.

Proof of (17.11). In the notation of (17.8), split the product A by the parity of the index. Using $\varepsilon_{2j} = \varepsilon_j$ and $\varepsilon_{2j+1} = -\varepsilon_j$,

$$A = \prod_{j \geq 0} \left(\frac{4j+1}{4j+2} \right)^{\varepsilon_j} \prod_{j \geq 0} \left(\frac{4j+3}{4j+4} \right)^{-\varepsilon_j} = \prod_{j \geq 0} \left(\frac{4j+1}{4j+3} \right)^{\varepsilon_j} \prod_{j \geq 0} \left(\frac{4j+4}{4j+2} \right)^{\varepsilon_j},$$

the second step regrouping the four linear factors. Since $(4j+4)/(4j+2) = (2j+2)/(2j+1)$, the last product is A^{-1} , so the quarter product equals $A^2 = 1/2$ by [Theorem 17.4](#): *the quarter product is the square of the Woods–Robbins product*. As with all such rearrangements, the manipulation is an exact identity at the level of even-length partial sums of the logarithmic series, and passes to the limit because both series converge. $\square$

Formal versions. The evaluation (17.11) is fully formal, with the proof above (`ThueMorseQuarterProduct.lean`): the parity split is the exact partial-sum identity “an even-length Woods–Robbins partial sum equals the quarter-ratio partial sum minus the Woods–Robbins partial sum” (`Fabius.wrA_two_mul`), whence $L(1/4, 3/4) = -\log 2$ (`Fabius.mpLimit_quarter`) and the product converges to $1/2$ (`Fabius.quarter_product`).

17.4 Three dyadic block-product limits

Dyadic truncation respects exact block recursion. The following limits are convenient fingerprints of the same product theory:

$$\lim_{m \rightarrow \infty} \prod_{k=0}^{2^m-1} \left(k + \frac{1}{2} \right)^{\varepsilon_k} = \frac{1}{2}, \quad (17.12)$$

$$\lim_{m \rightarrow \infty} \prod_{k=0}^{2^m-1} (k+1)^{\varepsilon_k} = \frac{1}{\sqrt{2}}, \quad (17.13)$$

$$\lim_{m \rightarrow \infty} \prod_{k=0}^{2^m-1} (k+1)^{(-1)^k \varepsilon_k} = \frac{1}{2\sqrt{2}}. \quad (17.14)$$

These limits are now *proved*, not merely recorded: over a dyadic block the even/odd split collapses the unbounded logarithms into the convergent sign-weighted ratio series. That split is isolated as a pair of hypothesis-free primitives over an arbitrary additive commutative group,

$$\sum_{k < 2N} \varepsilon_k g(k) = \sum_{j < N} \varepsilon_j (g(2j) - g(2j+1)), \quad \sum_{k < 2N} (-1)^k \varepsilon_k g(k) = \sum_{j < N} \varepsilon_j (g(2j) + g(2j+1))$$

(`Fabius.sum_zsmul_thueMorseSign_two_mul` and `Fabius.sum_zsmul_alternating_thueMorseSign_two_mul`, `ThueMorseEulerTransform.lean`); the alternating weight is exactly the one whose two halves add instead of cancelling, and the $\mathbb{Z}$ -module formulation makes the same two lines serve real logarithms, complex values, and formal coefficients alike. Applied to $g(k) = \log(k+1)$ they give the exact partial-sum identity $\sum_{k < 2N} \varepsilon_k \log(k+1) = L_A\text{-partial}(N)$ (`Fabius.block_log_one`), whence (17.13) follows from the Woods–Robbins limit (`Fabius.tendsto_block_product_one`); the alternating version telescopes into two Woods–Robbins partials — the $\log 2$ -contributions vanish by dyadic balance — giving (17.14) (`Fabius.tendsto_block_product_mixed`); and the half-shifted block sum is the master series at $(\frac{1}{4}, \frac{3}{4})$, so the limit in (17.12) is $\exp L(\frac{1}{4}, \frac{3}{4})$ (`Fabius.tendsto_block_product_half'`, `ThueMorseBlockProducts.lean`); its numerical value $1/2$ is the quarter product (17.11), proved above, and the evaluated limit is formal as well (`Fabius.tendsto_block_product_half`). Pairing adjacent factors reduces the first two to $G(1/4, 3/4)$ and $G(1/2, 1)$, respectively; the third factors into their product. The evaluations are thus (17.11) and (17.7).

The dyadic scaffolding is removable, and the shift is free. Nothing in the argument uses that the block length is a power of two, and nothing uses that the shift is 1 or $1/2$. For every real $c > 0$ and every M ,

$$\sum_{k < 2M} \varepsilon_k \log(k+c) = \sum_{j < M} \varepsilon_j \log \frac{j + \frac{c}{2}}{j + \frac{c+1}{2}}, \quad (17.15)$$

the order- M partial sum of $L(\frac{c}{2}, \frac{c+1}{2})$ (`Fabius.block_log_shift`): the pairing $2j \mapsto 2j+1$ uses only $2j+c = 2(j+\frac{c}{2})$ and $2j+1+c = 2(j+\frac{c+1}{2})$, so the two $\log 2$ terms cancel inside each pair and every unbounded term disappears at once. Exponentiating through `RealZPowProduct.lean` turns this into

$$\lim_{M \rightarrow \infty} \prod_{k=0}^{2M-1} (k+c)^{\varepsilon_k} = \exp L\left(\frac{c}{2}, \frac{c+1}{2}\right) \quad (c > 0), \quad (17.16)$$

a limit along *every* even block length, with no dyadic subsequence (`Fabius.tendsto_block_product_even`). The three fingerprints above are subsequences of (17.16) and of its alternating companion: (17.13) and (17.12) are the cases $c = 1$ and $c = 1/2$ at $M = 2^{m-1}$ (`Fabius.tendsto_block_product_one_even`, `Fabius.tendsto_block_product_half_even'`, and the evaluated `Fabius.tendsto_block_product_half_even`), while (17.14) is the case $M = 2^{m-2}$ of the mixed product along every multiple of four (`Fabius.block_log_mixed_four`, `Fabius.tendsto_block_product_mixed_four`). The three original statements survive verbatim, each a three-line subsequence corollary of its general form.

At $c = 1$ the general collapse (17.15) lands on $(\frac{1}{2}, 1)$, and there the master series is *termwise* the Woods–Robbins series, since $\log \frac{n+1/2}{n+1} = \log \frac{2n+1}{2n+2}$ (`Fabius.mpLog_one_half_eq_wra`). Hence

$$L(\tfrac{1}{2}, 1) = -\tfrac{1}{2} \log 2, \quad (17.17)$$

that is, $G(1/2, 1) = 1/\sqrt{2}$ (`Fabius.mpLimit_one_half_one`): [theorem 17.4](#) is the master log-series at $(\frac{1}{2}, 1)$, not merely a relative of it. Together with $L(1/4, 3/4) = -\log 2$ (`Fabius.mpLimit_quarter`) this evaluates (17.16) at both shifts the atlas uses.

18 Base- q digit characters and generalized Prouhet identities

Let $q \geq 2$, let $s_q(n)$ be the sum of the base- q digits of n , and let $\zeta \neq 1$ satisfy $\zeta^q = 1$. Define

$$u_n = \zeta^{s_q(n)}. \quad (18.1)$$

Theorem 18.1 (Generalized digit product). *For $0 \leq r < q$,*

$$u_{qn+r} = \zeta^r u_n. \quad (18.2)$$

For every $m \geq 0$,

$$\sum_{n=0}^{q^m-1} u_n z^n = \prod_{j=0}^{m-1} \left(\sum_{r=0}^{q-1} \zeta^r z^{r q^j} \right). \quad (18.3)$$

Proof. Appending the base- q digit r increases the digit sum by r . In the product, choose one digit independently at each position; the exponent is the corresponding base- q integer and the coefficient is ζ raised to its digit sum. $\square$

Formal versions. [Theorem 18.1](#) is formal in `BaseDigitProduct.lean`, and in a form that needs no root of unity at all. The recursion (18.2) is `Fabius.zeta_pow_digits_sum_mul_add`, resting on the digit-sum law $s_q(qn + r) = r + s_q(n)$ for $r < q$ (`Fabius.digits_sum_mul_add`); the product (18.3) is `Fabius.sum_zeta_pow_digits_sum_mul_pow`. Both hold for *every* ζ in a commutative semiring: the hypothesis $\zeta^q = 1$ is never used, because the identity is a statement about assembling an integer from independently chosen digits, not about ζ . The formal proof makes this explicit by first proving the product for an *arbitrary* digit weight $w : \mathbb{N}_0 \rightarrow R$,

$$\sum_{n < q^m} \left(\prod_{j < m} w(d_j(n)) \right) z^n = \prod_{j < m} \left(\sum_{r < q} w(r) z^{r q^j} \right),$$

where $d_j(n) = \lfloor n/q^j \rfloor \bmod q$ is the j -th base- q digit (`Fabius.sum_windowWeight_mul_pow`, with `Fabius.digitAt` and `Fabius.windowWeight`); the theorem in the text is the case $w(r) = \zeta^r$, and the bridge from the digit window to $s_q(n)$ is `Fabius.sum_digitAt_eq_digits_sum`.

Each factor vanishes at $z = 1$, since $1 + \zeta + \dots + \zeta^{q-1} = 0$ (`Fabius.geom_sum_eq_zero_of_pow_eq_one` in an integral domain, from $\zeta^q = 1$, $\zeta \neq 1$).

Theorem 18.2 (Generalized Prouhet cancellation and sharp moment). *If p is a polynomial of degree less than m , then*

$$\sum_{n=0}^{q^m-1} \zeta^{s_q(n)} p(n) = 0. \quad (18.4)$$

At the first surviving degree,

$$\sum_{n=0}^{q^m-1} \zeta^{s_q(n)} n^m = m! \frac{q^{m(m+1)/2}}{(\zeta - 1)^m}. \quad (18.5)$$

Proof. Set $z = e^t$ in (18.3). Every factor has a zero at $t = 0$, proving cancellation below degree m . Its derivative at zero is

$$q^j \sum_{r=0}^{q-1} r \zeta^r = q^j \frac{q}{\zeta - 1} = \frac{q^{j+1}}{\zeta - 1}, \quad (18.6)$$

where the finite geometric sum was differentiated. Multiplying the m linear terms yields $q^{1+2+\dots+m}/(\zeta - 1)^m$; coefficient extraction contributes $m!$. $\square$

Formal versions. Both displays are formal in `BaseDigitProuhet.lean`, with $P_d(m) = \sum_{n < q^m} \zeta^{s_q(n)} n^d$ as `Fabius.digitPowerSum`. The cancellation (18.4) is `Fabius.digitPowerSum_eq_zero_of_lt` for monomials, $P_d(m) = 0$ for $d < m$, and `Fabius.sum_zeta_pow_digits_sum_mul_eval_eq_zero` for an arbitrary polynomial of degree below m ; both take the single hypothesis $\sum_{r < q} \zeta^r = 0$ over any commutative semiring, and $\zeta^q = 1$ is not used. The sharp moment

(18.5) is proved division-free, as $(\zeta - 1)^m P_m(m) = m! q^{0+1+\dots+m}$ in any commutative ring (Fabius.us.sub_one_pow_mul_digitPowerSum_self), and then in the displayed form with $/(\zeta - 1)^m$ and the Gauss sum evaluated, in a field with $\zeta \neq 1$ (Fabius.digitPowerSum_self_eq_div). The engine is the one-step recursion Fabius.digitPowerSum_succ, obtained by splitting $[0, q^{m+1})$ into q blocks $q^m r + a$, where the digit sum adds r (Fabius.digits_sum_pow_mul_add) and the binomial theorem expands $(a + q^m r)^d$; in that expansion every $P_i(m)$ with $i < d$ vanishes by induction, so for $d \leq m$ only $i = d$ survives with coefficient $\sum_r \zeta^r = 0$, while at $d = m + 1$ the survivor $i = m$ carries $(m + 1)q^m \sum_r r \zeta^r$. The root-of-unity derivative (18.6) is the telescoping identity Fabius.geom_sum_mul_weighted, $(\zeta - 1) \sum_{r < q} r \zeta^r = (q - 1)\zeta^q + 1 - \sum_{r < q} \zeta^r$ for every ζ , specialized in Fabius.sub_one_mul_sum_mul_pow. The degree-zero case is also available on its own as Fabius.sum_zeta_pow_digits_sum_eq_zero in BaseDigitProduct.lean, by evaluating the digit product at $z = 1$.

For $q = 2$ and $\zeta = -1$, the sequence is $u_n = \varepsilon_n$ and (18.5) becomes $(-1)^m m! 2^{\binom{m}{2}}$, exactly (9.24). The binary theory is therefore the first member of a root-of-unity digit-character family, while the particularly rich valuation, carry, and Pascal-row formulas remain genuinely base-two features.

19 Formula selection, computational checks, and a Lean roadmap

19.1 Formalized modules (August 2026)

The first tranches of the formalization now exist as one hundred and three compiled Lean modules in Analysis/FabiusFunction, complementing the corpus modules (Parity.lean, TwoAdic.lean, ThueMorsePrefix.lean, ThueMorseGenerating.lean, ThueMorseExponential.lean, ThueMorseBinomialLog.lean) that already covered the recursion, XOR, reflection, Lucas/Glaisher, binomial–logarithm, modulo-three formula, generating products, Prouhet moments, and prefix theory. The count is exactly the number of rows of the table below, one row per module; the proposal table of Section 19.4 re-lists some of these same modules under the boundaries that were originally proposed for them, and is a roadmap rather than a second inventory. The series-filter and fractional-Volterra rows were checked at source checkpoint 149332f9d.

Module	Contents
ThueMorseBitSupport.lean	Canonical binary support and two-power equivalence, both reconstruction directions, test-bit and empty/zero laws, shifts, weight and sign laws, support bounds, the bounded Boolean-cube equivalence including level zero, image equality, and powerset-sum reindexing.
ThueMorseBooleanCube.lean	Two-parameter master product over any commutative semiring, signed product over any commutative ring, and weight enumerator, using the centralized bit-support reindexing API.
ThueMorseComplexExponential.lean	Finite complex exponential transforms of dyadic blocks: Hamming-weighted base and arbitrary affine forms, and signed Thue–Morse base and arbitrary affine forms, all including level zero. The separate ThueMorseComplexProductBridge.lean supplies the total complex sinc and negative-Laplace normalizations.
ThueMorseValuation.lean	Successor law, weight subadditivity, ruler product, Catalan valuation, additive Kummer, borrow cocycle.
CentralBinomialValuation.lean	Exhaustive three-theorem bridge: central-binomial valuation equals binary weight, its parity gives the Thue–Morse sign, and zero valuation is equivalent to zero binary weight (hence occurs only at $n = 0$).

Module	Contents
<code>BinaryDigitFloor.lean</code>	Exhaustive four-theorem total-natural bridge from dyadic quotients to floor-difference binary digits: quotient-by-two normalization, parity remainder, the consecutive-floor difference, and the test-bit indicator identity; no analytic or real-floor generalization.
<code>ThueMorseDigits.lean</code>	Additive Legendre, floor sums, floor-parity formula, ‘testBit’ bridge, digit product.
<code>RademacherSine.lean</code>	Exhaustive nine-theorem square-wave and Rademacher–sine closure: the everywhere-valid factored sine identity; positivity and sign under exactly nonzero fractional part; the half-shifted floor, its unconditional nonintegrality, the exact bit factor, and eventual factor $+1$ under $n < 2^j$; and the finite product under $n < 2^m$ plus the total $\prod'$ identity, reduced to finite support rather than analytic infinite-product convergence.
<code>ThueMorseEnumerators.lean</code>	Weight distribution, evil/odious enumerations and composition laws, XOR autocorrelation, signed Hamming spheres.
<code>FinitePolynomialFunctional.lean</code>	Moment expansion for finite weighted polynomial evaluations, selected-, top-, and constant-coefficient extraction, scalar-extension forms, degree-valued strict cancellation, and same-ring conveniences over commutative semirings; repeated nodes and a zero surviving moment are allowed.
<code>FinitePowerSeriesFilter.lean</code>	The universal finite rescaling filter on formal power series over arbitrary commutative semirings: diagonal coefficient action by the finite node moments, constant-term preservation from mass one or exactness, and compatibility with further rescaling and coefficient maps. This is coefficientwise algebra only; it makes no analytic evaluation or convergence claim.
<code>CompleteHomogeneous.lean</code>	Evaluated complete homogeneous symmetric polynomials over arbitrary commutative semirings: explicit multiset and universal-symmetric-polynomial forms, adjoining-variable and map/reindex laws, and the univariate complete-homogeneous quotient; only the linear-factor identity requires subtraction.
<code>CompleteHomogeneousGenerating.lean</code>	Finite formal generating series for complete homogeneous evaluations: exact coefficients, empty-family and adjoining-variable recurrences over commutative semirings, and one-factor/full-product denominator clearing plus the canonical <code>PowerSeries.invOfUnit</code> reciprocal over commutative rings. Repeated values, zero divisors, and nilpotents are allowed; analytic evaluation, convergence, and the Bell/generalized-harmonic conversion are not asserted.
<code>LagrangeResidualMoments.lean</code>	The all-order residual-moment formula for any finite weighted row exact below the node count, over an arbitrary commutative ring with repeated nodes allowed; distinct field-valued nodes enter only in the Lagrange-weight specialization.
<code>BaseLayerRegrouping.lean</code>	Finite base- b layer reindexing and regrouping over arbitrary commutative monoids, with additive companions; the truncated divisor count and its evaluation as $1 + \nu_b(m)$ under explicit layer-range coverage. No infinite-product or convergence assertion is made.
<code>GeometricCompleteHomogeneous.lean</code>	The denominator-free principal specialization $h_n(1, q, \dots, q^r) = \left[\begin{smallmatrix} n+r \\ r \end{smallmatrix} \right]_q$ over arbitrary commutative semirings, including arbitrary common scale and the range-indexed form.

Module	Contents
<code>GeometricResidualMoments.lean</code>	Every positive moment of an arbitrary exact geometric row over any commutative ring, with repeated nodes and arbitrary scales allowed; geometric Lagrange rows, common dilations, and shifted blocks follow over fields under finite-node injectivity.
<code>GeometricPowerSeriesFilter.lean</code>	Formal-series filtering at the nodes $1, q, \dots, q^p$: coefficient cancellation through the exactness order, the Gaussian multiplier in every positive degree, an all-index residual-series identity with the mass-one degree-zero boundary, and the geometric Lagrange specialization. Analytic substitution, convergence, and remainder bounds are not asserted.
<code>AnalyticSeriesFilter.lean</code>	Convergent-series form of an arbitrary finite rescaling filter: diagonal action, exact finite head and infinite tail, the target-zero constant boundary, and geometric and geometric-Lagrange Gaussian tails, including shifted scales. <code>Mathlib Summable</code> is unconditional; hypotheses are imposed on weighted samples, so a zero-weight node needs no unweighted convergence. Conditionally-only boundary series, signs, bounds, uniform convergence, and sinc-product instantiations are not asserted by this generic module.
<code>RvachevQBinomialFilter.lean</code>	Exact specialization of the Gaussian filter to the actual infinite Rvachev product: the named even Fourier modes sum to the product and scale by powers of c^2 ; the generic complex theorem assumes exactly injectivity of the finite powers of c^2 , while $c = 1/2, q = 1/4$ is assumption-free. This does not identify the finite prefixes Q_m or $P_{b,N}$, their quotient or Bell coefficients, or prove the report’s analytic acceleration estimates.
<code>FractionalVolterraCalculus.lean</code>	Positive affine covariance on ordered real intervals for arbitrary real order and Gamma-normalized order raising for real $\alpha > 0$ in a complete normed real space, assuming continuity of the primitive, interval integrability of its displayed derivative, and the right derivative <code>HasDerivWithinAt</code> on the open interval; includes the zero-boundary form. For $\alpha \leq 0$, affine covariance is an identity for the totalized Lean interval-integral definition; a classical Riemann–Liouville or integrability interpretation is asserted only for positive order.
<code>FabiusFractionalVolterra.lean</code>	Total causal Rvachev fractional primitive with its support-truncated integral, positive-natural bridge, and positive-order semigroup on $x \geq -1$, plus boundary-free order raising and affine covariance specialized to the dyadic system: the signed extension for $x \geq 0$, bounded Fabius for $0 \leq x \leq 1$, and the up-to-Fabius bridge based at -1 for $x \geq -1$, all for real $\alpha > 0$. Complex, weighted-monomial, iterated, negative-branch, shifted-lattice, endpoint-moment, transform/tail, piecewise/refinement, and inverse/quantile extensions remain outside the module.
<code>LogSeriesMultiplicative.lean</code>	The recurrence-defined formal power-series logarithm over arbitrary commutative $\mathbb{Q}$ -algebras: coefficient reconstruction, zero constant term, differential-equation characterization and uniqueness, binary product rule, unit law, and finite-product additivity; no analytic-logarithm or infinite-product claim.

Module	Contents
<code>FiniteQBinomialCore.lean</code>	Denominator-free Gaussian coefficients by both q -Pascal recurrences, functoriality under semiring maps, and reflection symmetry over commutative semirings; the integer-index extension <code>gaussianBinomialInt</code> , with its natural-index, negative-index, out-of-range, and reflection laws; finite q -Pochhammer shift/concatenation laws, the full product–Gaussian factorial identities, and the finite q -binomial theorem over commutative rings, including zero and roots of unity.
<code>QBinomialVandermonde.lean</code>	Both zero-extended q -Vandermonde orientations, the full-range and exact-min-support central identities, the total positive shifted-central form, and two negative-shift forms under exactly $k \leq N$, plus the total integer-shift identity <code>gaussianBinomial_two_mul_int_shifted_central</code> and its finitely supported <code>finsum</code> form, all denominator-free over arbitrary commutative semirings.
<code>GeometricQBinomialLagrange.lean</code>	The reversed finite q -binomial product for geometric Lagrange numerators, their global forward-index Gaussian form via division-free symmetry and zero extension, normalized weights under the intrinsic node-injectivity hypothesis, and every positive residual moment.
<code>ThueMorseSparseProuhet.lean</code>	Arbitrary-weight Prouhet annihilation and exact top-moment formula (theorem 9.4), degree-valued wrappers covering the zero polynomial on empty supports, the direct affine block extraction and cancellation API (theorem 9.6), and all-index sparse submask identities.
<code>ThueMorseRarefied.lean</code>	Modulo-three rarefaction recursion and closed vectors, Newman positivity at dyadic endpoints, block balance, even-modulus collapse.
<code>ThueMorseAperiodicity.lean</code>	Block-flip identities, promotion of an eventual period to a global one, and non-eventual-periodicity — the last now deduced from overlap-freeness rather than from the reflection argument.
<code>ThueMorseBlockAlgebra.lean</code>	Telescope and multiplicity ladder, degree and leading data of P_m , reciprocity (reverse and evaluation forms), cyclotomic multiplicity table, binomial-basis vanishing and sharp moment over any commutative ring.
<code>ThueMorseMoments.lean</code>	Hockey-stick column sum, complete binomial moment composition (corollary 9.14), reflection functional equation, midpoint parity selection rule, sharp centered moment; the master products on arbitrary supports (including the arbitrary-exponent form), the exponential substitution, the complete sparse moment composition on an arbitrary support, and the complete power-moment composition formula as its dyadic-block case.
<code>ThueMorseArtinSchreier.lean</code>	Characteristic-two power-series Frobenius, the $\mathbb{F}_2$ bit series with its functional equation, Artin–Schreier quadratic and normalized forms, the complete two-element solution set $\{\Theta, \Theta + G\}$ with zero-constant-term uniqueness as its corollary, finite Artin–Schreier telescope.
<code>ThueMorseEulerTransform.lean</code>	Prefix collapse and ones-count formulas, general residue-class reindexing, ruler-convolution recurrence by finite combinatorics, interval discrepancy bounds.
<code>ThueMorseLucasSupport.lean</code>	Lucas’s submask criterion (bit, <code>&&&</code> , and support forms), odd Pascal positions as the submask cube, signed Pascal-support product, balance on odd positions.

Module	Contents
<code>ThueMorseWalsh.lean</code>	All-ones digit patterns and masked two-power sums, ε as a Walsh character, character orthogonality on the block, the one-point Walsh spectrum.
<code>ThueMorseSparseMoments.lean</code>	The atlas’s all-sparse-moments formula on the submask cube (corollary 10.5), as the $S = J(n)$ instance of the arbitrary-support composition of <code>ThueMorseMoments.lean</code> .
<code>ThueMorseMixedDifference.lean</code>	The dyadic mixed-difference operator and the finite-difference factorization for arbitrary functions from an additive commutative monoid into an additive commutative group, with the powerset expansion.
<code>ThueMorseSymmetricDifference.lean</code>	Exhaustive two-definition/eleven-theorem centered companion: arbitrary half-steps and the symmetric Boolean cube, sharp commutative-ring polynomial extraction/cancellation/top value, the centered dyadic block, and both characteristic-zero increasing-grid forms. It does not supply the repeated-integral clause of the continuous-chaos report.
<code>ThueMorseAutocorrelation.lean</code>	Finite autocorrelation recursion in exact integer form along <i>any even</i> window, $R_{2K}(2r) = 2R_K(r)$ and $R_{2K}(2r+1) = -(R_K(r) + R_K(r+1))$, with the dyadic layer $A_m(k) = R_{2^m}(k)$ as its power-of-two case: zero-shift value, trivial bound, closed values at shifts one and two, and the normalized limits at shifts zero, one and two.
<code>ThueMorseOverlapFree.lean</code>	Thue’s overlap-freeness theorem with cube-freeness, the attained critical exponent two, the alphabet-free implication “overlap-free $\Rightarrow$ not eventually periodic”, and constructive complement and reversal closure of the factor language.
<code>ThueMorseRecurrence.lean</code>	Full block-concatenation laws for weight and bit, all two-letter factors in every fourteen-letter window, uniform recurrence with an explicit linear bound.
<code>ThueMorseMahler.lean</code>	Mahler equation and dyadic dissection for the signed series, cleared zero–one bridge, prefix generating function, finite-level autocorrelation Mahler equation.
<code>ThueMorseBooleanMobius.lean</code>	Submask complement laws, submask weight enumerators, reconstruction and inversion-polynomial identities, nested binomial formula, Boolean Möbius inversion on finsets and on submasks.
<code>ThueMorseMultinomial.lean</code>	Kummer’s criterion for multinomial coefficients, the multinomial support count $O_r(n) = r^{\text{wt}_2(n)}$, the signed r -nomial row formula, the exact r -nomial logarithm.
<code>ThueMorseFourier.lean</code>	The circle factor $ 1 - e^{i\theta} ^2 = 2 - 2\cos\theta$ and the finite Riesz product with its nonnegativity.
<code>ThueMorseSineProduct.lean</code>	The half-angle factorization, the complex sine-product closed form of the Thue–Morse trigonometric polynomial, and its arbitrary-phase affine real sine form.
<code>ThueMorseComplexProductBridge.lean</code>	Total complex finite Thue–Morse sinc and negative-Laplace bridges at every level: complex-argument sine and half-angle forms, normalized prefixes with empty-level, successor, and origin laws, denominator-free block equalities valid at the removable origin, quotient normalizations away from zero, and the exact finite Fourier–Laplace rotation.
<code>UniformDigitThueMorseBridge.lean</code>	The exact finite probability consumer: the characteristic function of the centered uniform-digit prefix is the normalized sinc prefix at every level and real frequency, hence satisfies the total Thue–Morse block identity; the reverse quotient formula assumes nonzero frequency.

Module	Contents
<code>LacunaryProductToSum.lean</code>	Binary-weight complementation on a dyadic block, the induced global-sign reflection, quarter-turn cosine parity, and the exact real lacunary sine product as a Thue–Morse sum.
<code>ThueMorseFourierInversion.lean</code>	Root-of-unity orthogonality and finite Fourier inversion over any field with a primitive root, specialized to the dyadic DFT over $\mathbb{C}$.
<code>SummableCyclicAlias.lean</code>	Finite-range scalar twists of summable families in complete uniform additive groups, and bilateral integer residue sums with forward and inverse $\mathbb{Z}\text{Mod}$ Fourier transforms in complete Hausdorff uniform complex modules; no norm or absolute-convergence hypothesis is imposed.
<code>ThueMorseParseval.lean</code>	Character orthonormality on $[0, 1]$ and the Parseval mass $\int_0^1 P_m(e^{2\pi i x}) ^2 dx = 2^m$.
<code>ThueMorseComplexity.lean</code>	Factor complexity: the supertile-pair window parametrization with the boundary-crossing case split, uniform-recurrence finiteness, and the sharpened two-sided bound $\ell + 1 \leq p(\ell) \leq 6\ell - 4$ – with $p(\ell) \leq 2 \cdot 2^k + 2(\ell - 1)$ at every covering level and the older $4 \cdot 2^k, p(\ell) < 8\ell$ kept as corollaries – and the Morse–Hedlund strict-growth argument.
<code>ThueMorseArithmeticFilter.lean</code>	The root-of-unity filter for arithmetic progressions over any field (division-free form), the boxed rarefied formula, and the even-modulus collapse.
<code>ToeplitzHessenberg.lean</code>	Cramer’s rule for a known solution over any commutative ring; the triangular and the lower Toeplitz–Hessenberg band matrices of arbitrary band data; the triangular determinant $\prod_{i < n} d_i$; and the expansion-free Toeplitz–Hessenberg determinant formula, read off from any solution of the convolution system, in matrix and in scalar-recurrence form.
<code>ThueMorseHessenberg.lean</code>	The ruler convolution as an instance of that band system: the triangular ruler system with determinant $n!$, the Hessenberg formula $\det H_n = (-1)^n n! \varepsilon_n$ for every $n \geq 0$ by one column rotation, and the sharp absolute value $ \det H_n = n!$.
<code>ExponentialPartition.lean</code>	Weighted multiplicity vectors and reciprocal-factorial partition sums over every commutative $\mathbb{Q}$ -algebra; the marked-part and normalized successor recurrences; and the quotient-notation bridge for characteristic-zero fields.
<code>ThueMorsePartition.lean</code>	Ruler-recurrence uniqueness and the partition and explicit Bell-polynomial formulae for ε_n , obtained by specializing the generic partition engine.
<code>ExponentialBell.lean</code>	Equality of the finite weighted-partition coefficient function with <code>SaddleExpansion.expCoeff</code> over every commutative $\mathbb{Q}$ -algebra, by recurrence uniqueness.
<code>MomentCumulantAlgebra.lean</code>	Factorial normalization and denormalization; the complete Bell and inverse moment–cumulant transforms; their classical recurrence, normalized all-index inverse laws, and functoriality over every commutative $\mathbb{Q}$ -algebra.
<code>PolynomialExpectationCumulant.lean</code>	Finite polynomial integrals from support-local raw moments, their complete-Bell/formal-cumulant rewrite with the exact zeroth-mass correction for arbitrary measures, and the normalized and probability-measure specializations.
<code>ThueMorseIntegerLift.lean</code>	The integral algebraic lift constructed in closed Catalan–binomial form; Catalan parity at Mersenne indices; the unconditional parity theorem $c_n \equiv \tau(n), \varepsilon_n = (-1)^{c_n}$.

Module	Contents
<code>ThueMorseBoundaryFlatness.lean</code>	The lacunary product $\mathcal{E}(t)$ as an infinite product: self-similarity, positivity, monotonicity, the effective upper bound $\mathcal{E}(t) \leq 2^{\binom{m}{2}} t^m$ at every level m , and the matching lower bound $2^{\binom{m}{2}} (t/2)^m \mathcal{E}(\frac{1}{2}) \leq \mathcal{E}(t)$ under the dyadic matching $\frac{1}{2} \leq 2^m t \leq 2$.
<code>ThueMorseDirichlet.lean</code>	The shifted Dirichlet series in the real regime: conditional convergence for every $\sigma > 0$ by Dirichlet’s test, and the dyadic parameter equation $D(\sigma, a) = 2^{-\sigma} (D(\sigma, \frac{a}{2}) - D(\sigma, \frac{a+1}{2}))$.
<code>ThueMorseCatalanSeries.lean</code>	The diagonal Vandermonde identity $\sum_i C(i, p) C(M-i, q) = C(M+1, p+q+1)$ and the second-difference bridge from the lift to the substituted Catalan series.
<code>ThueMorseCatalanConvolution.lean</code>	The interior Catalan convolution and the quadratic identity $\sum_{p+q=m} h_p h_q = h_m - h_{m+1} - [m=0]$ of the substituted series.
<code>ThueMorseIntegerLiftEquation.lean</code>	The boxed algebraic equation $(1-z)^3 C^2 + (1-z)^2 C = z$ in $\mathbb{Z}[[z]]$ for the constructed lift, by identification with the substituted series and cancellation.
<code>ThueMorseIntegerLiftRecurrence.lean</code>	The linear ODE $(1+3z)(1-z)^2 C' = 1 + 3(1-z^2)C$ by differentiation and one linearization certificate, and the three-term recurrence by telescoping its coefficients.
<code>ThueMorseLacunaryQuotient.lean</code>	The dyadic geometric series and its inversion of $1 - z^{2^j}$; the denominator-free bit Mahler equation; the lacunary quotient expansion at every finite level.
<code>ThueMorseInfiniteProduct.lean</code>	The analytic infinite-product identity $\sum \varepsilon_n z^n = \prod_j (1 - z^{2^j})$ on the complex unit disc $\ z\ < 1$, including absolute convergence and multipliability; the real interval $ x < 1$ and exponential ray $\sum \varepsilon_n e^{-nt} = \mathcal{E}(t)$, $t > 0$, are retained as specializations.
<code>ThueMorseMellin.lean</code>	The Mellin representation $\Gamma(\sigma) D(\sigma, a) = \int_0^\infty t^{\sigma-1} e^{-at} \mathcal{E}(t) dt$ for $\sigma > 1$: kernel integrability, absolute interchange, and termwise Gamma evaluation.
<code>RealZPowProduct.lean</code>	Finite products of integer powers of positive reals through the exponential: the signed-logarithm identity, the limit transfer from logarithm sums to products along an arbitrary filter, and the closed-form value computations; nothing in it mentions Thue–Morse.
<code>ThueMorseWoodsRobbins.lean</code>	The Woods–Robbins evaluation $\prod ((2n+1)/(2n+2))^{\varepsilon_n} = 2^{-1/2}$: Dirichlet-test convergence of the three log-series, the termwise and interleaving identities, and exponentiation.
<code>ThueMorseMasterProduct.lean</code>	The master product $G(a, b)$: convergence for all positive parameters, the cocycle identities, the dyadic functional equation, and $G = \exp L$.
<code>ThueMorseBlockProducts.lean</code>	Block-product limits along every even block length and for an arbitrary positive shift c : the shift-general log collapse onto the master partials at $(\frac{c}{2}, \frac{c+1}{2})$, the identification of Woods–Robbins as the master series at $(\frac{1}{2}, 1)$, the alternating telescoping along every multiple of four, and the three dyadic limits as subsequence corollaries.
<code>ThueMorseEntireContinuation.lean</code>	The entire continuation: flatness of every order at zero, exponential decay, local integrability, entirety of $\Gamma^{-1} \cdot \text{Mellin}$, the trivial zeros at the nonpositive integers, agreement with $D(\sigma, a)$ on $\sigma > 1$, and the dyadic parameter equation at every complex s by the identity theorem.

Module	Contents
<code>FrullaniIntegral.lean</code>	The Frullani integral $\int_0^\infty (e^{-pt} - e^{-qt})t^{-1} dt = \log(q/p)$, by differentiation under the integral sign; not in Mathlib.
<code>ThueMorseGDirichlet.lean</code>	The derivative bridge $\log G(a, b) = D'(0, b) - D'(0, a)$: the closed integral form of $\log G$, $(1/\Gamma)'(0) = 1$, and the derivative of the entire continuation at zero.
<code>ReciprocalGammaJets.lean</code>	Exact first jets of the entire reciprocal Gamma function at every nonpositive integer: derivative $(-1)^r r!$, simple analytic order, the punctured local coefficient, and the retained origin compatibility theorem; no derivative is assigned to raw Gamma at a pole.
<code>ThueMorseGammaTower.lean</code>	Exact first jets of the entire Dirichlet–Mellin continuation, total GammaLog and exponentiated tower definitions, Mellin and real-integral GammaLog formulas, dyadic log and tower laws for positive parameter, and the positive-parameter level-zero master-product ratio; GammaLog is a chosen coordinate, while the parameter-differential ladder remains outside the module.
<code>ThueMorseQuarterProduct.lean</code>	The quarter product $G(1/4, 3/4) = 1/2$ from Woods–Robbins by parity splitting, and the value $1/2$ of the half-shifted block product.
<code>ThueMorseDiscSeries.lean</code>	The Thue–Morse power series as a holomorphic function on the unit disc: the complex-disc product wrapper $F(z) = \prod_{j=0}^\infty (1 - z^{2^j})$, locally uniform convergence, the analytic Mahler equation and its iterate, and the real-ray bound $ F(x) \leq 1 - x$.
<code>ThueMorseNaturalBoundary.lean</code>	The natural-boundary theorem: radial vanishing at dyadic roots of unity, two-power roots approximating every circle point, and the double identity-theorem argument.
<code>ThueMorseComputability.lean</code>	Primitive recursiveness: the descent-recursion scheme for course-of-values recursions, the halving identity for the binary weight, and the primitive recursive and computable forms of wt_2 , τ , and ε .
<code>ThueMorseNewman.lean</code>	Newman’s positivity theorem: residue-class sign sums, the modulo-three parity system, the base-four collapse with disjoint boundary corrections, and the strong induction.
<code>ThueMorseGelfond.lean</code>	Gelfond’s sharp uniform bound: the one-variable sum-of-squares inequality $\varphi^2(t)\varphi(2t) \leq 3\sqrt{3}$, the telescoped doubling-orbit product bound, and the exponential-sum form with constant $2/\sqrt{3}$.
<code>PeriodDoublingHankel.lean</code>	The mod-two Hankel engine: reusable block-determinant eliminations, parity interleavings, the period-doubling bit, and the four determinant halvings forcing every $\mathbb{F}_2$ -determinant to equal one.
<code>ThueMorseHankel.lean</code>	The Allouche–Peyrière–Wen–Wen theorem: the doubling factorizations $H_{2n} = 2^n H_n H_n(m)$, $H_{2n+1} = 2^n H_{n+1} H_n(m)$, oddness of the difference determinants, and the exact valuation $\nu_2(H_r) = r - 1$.
<code>LacunaryRieszIntegral.lean</code>	The exact orthogonality theory of lacunary Riesz products, with no Thue–Morse content: the frequency detector, the master theorem $\int_0^1 \cos(2\pi Kt + \psi) \prod_{j < n} (1 + a_j \cos(2\pi m_j t + \varphi_j)) dt$ under the dissociation hypothesis $ K + \sum_{i < j} m_i < m_j$, its super-increasing and Hadamard-gap corollaries, the mean value 1, and the exact identity $\int_0^1 \prod_{j < n} \sin^2(\pi m_j t) dt = 2^{-n}$ with its geometric and dyadic specializations.

Module	Contents
<code>GelfondLogisticBound.lean</code>	The Gelfond rate from the logistic map, again with no Thue–Morse content: invariance of $[0, 1]$, the potential telescope giving $\prod_{j < n} u_j \leq \frac{5}{3}(\frac{3}{4})^n$ along any orbit from an initial value in $[0, 1]$, exactness of the rate at the fixed point $3/4$, the squared and unsquared dyadic sine-product bounds, and the L^1 lower bracket for the sine product.
<code>VarianceBookkeeping.lean</code>	Closed evaluation of the stationary covariance sum, including the zero-length variance identity.
<code>BirkhoffVariance.lean</code>	Exact pair correlations for the logarithmic doubling cocycle and its all-index Birkhoff square mean.
<code>BirkhoffProductBridge.lean</code>	The almost-everywhere bridge from the cocycle to the normalized dyadic sine product and the all-index logarithmic product variance (3.20).
<code>FourierProduct.lean</code>	The Rvachev sinc product and its identification with the Fourier transform, including the exact complex zero set $\Phi(z) = 0 \iff z \in \mathbb{Z} \setminus \{0\}$.
<code>RenormalizationIdentity.lean</code>	The standalone product identity $\Phi(z) = \text{sinc}_{\text{rad}}(\pi z)\Phi(z/2)$ and its exact n -fold shell form.
<code>IntegerZeroLocalFactorization.lean</code>	Total denominator-clearing identities for shifted and arbitrarily complex-scaled sinc factors; the global local-coordinate factorization at $2^k q$, its odd-center form, and canonical positive-natural and signed-integer forms using the exact dyadic valuation and odd part.
<code>LeadingJet.lean</code>	Vector-valued, field-generic computation of every jet of $z^n \cdot U(z)$, including the all-index zero-or-binomial-factorial formula, division-free $n + r$ forms, scalar specializations, and germ versions at zero and arbitrary centers; only the corresponding finite differentiability of U is required.
<code>IntegerZeroAnalyticOrder.lean</code>	Entirety of the standalone sinc product; nonvanishing at odd half-integers; the total integer-zero local cofactor, its shifted and centered germ factorizations, analyticity away from the totalized denominator point and exact central value; exact complex analytic order at every nonzero integer; all lower-jet vanishings, every higher cofactor jet, and the exact nonzero leading derivative.
<code>SincProductShells.lean</code>	Exact dyadic shell factorization of $ \Phi $, its total cross-multiplied form at every real seed (including zero and the empty shell), and its Gelfond-rate bound along every nonzero real dyadic ray.
<code>SincCanonicalProduct.lean</code>	The canonical product $\prod_{m \geq 1} (1 - z^2/m^2)^{1+\nu_2(m)}$, obtained by regrouping the absolutely convergent double Euler product.
<code>PhiZeroOrder.lean</code>	The exact positive-integer real-axis zero order $1 + \nu_2(m)$, expressed by a punctured norm quotient converging to a positive constant.
<code>ThueMorseLobeSign.lean</code>	The lobe identity $\Phi(x) = \varepsilon_m \Phi(x) $ for $m < x < m + 1$, from the closed lattice-point count.
<code>RvachevProductContinuity.lean</code>	Real-axis continuity of Φ and $ \Phi $; the origin as the unique unit-modulus point, with strict inequality elsewhere; and existence and uniqueness of the modulus maximum on every positive lobe $(m, m + 1)$ and reflected negative lobe $(-(m + 1), -m)$, $m \geq 1$.
<code>RvachevFourierDecay.lean</code>	Public facade transferring the canonical product, total and divided shell formulas, global decay, strict peak, extremal ray, lobe signs and maxima, and zero orders from Φ to the actual Rvachev Fourier transform.

Wherever the Lean statement is stronger than the prose — arbitrary weights, arbitrary coefficient rings, arbitrary moduli and supports — the text above records the stronger form. The original proposal table further below is retained with its planned module boundaries for comparison; most rows have since been delivered under the module names above, as marked row by row.

19.2 Which representation is best?

The formulas are equivalent mathematically but not proof-theoretically or computationally. The following table records the most economical interface for common tasks.

Goal	Recommended representation	Reason
Executable evaluation	Recurrence or machine popcount	Linear in the bit length; no large intermediate integers.
Dyadic reindexing	Block concatenation (2.6)	Quotient and remainder separate without carries.
Bitwise algebra	XOR character (2.9)	Makes ε a genuine group character.
Ordinary addition	Carry cocycle (4.10)	Isolates the exact obstruction to XOR multiplicativity.
Number-theoretic disguise	Central binomial valuation (4.4)	One classical integer stores the full binary weight.
Pascal-row interpretation	Glaisher count or modulo-three residue	Uses the odd support of one row; the residue form avoids logarithms.
Higher-dimensional Pascal	Multinomial count (5.18)	Works with any number of variables.
Formal and convergent power series	Formal products plus coefficient-diagonal and analytic filters	Coefficients and Mahler equations remain algebraic; exact analytic filtering is available under unconditional summability, with zero-weight nodes discharged sharply, and the named even modes now instantiate the actual infinite Rvachev product. Conditional boundary convergence, the finite-prefix Q_m or $P_{b,N}$ quotient, signs, bounds, and uniform control remain separate obligations.
Nonlocal coefficient formulas	Euler/Bell/determinant forms	Eliminate explicit digit notation and expose ruler convolutions.
Finite polynomial row	Moment-functional extraction	Separates polynomial coefficients from row moments, supports scalar extension, and does not require distinct nodes or a nonzero surviving moment.
Polynomial cancellation	Mixed finite differences	Proves all Prouhet moments in one operator identity.
Arbitrary Pascal row	Sparse signed product (10.2)	Restricts finite differences to the odd row support.

Goal	Recommended representation	Reason
Finite harmonic analysis	Walsh character	The entire block is one pure frequency.
Ordinary spectral analysis	DFT sine product and Riesz density	Exposes roots of unity, amplitudes, and correlations.
Analytic continuation	Mellin integral (16.11)	Boundary flatness removes every Mellin singularity.
Infinite products	Master product $G(a, b)$	Pair reduction and cocycle laws organize evaluations.
Fabius approximation	Iterated-prefix generating function	Coefficient grids arise from $E(z)/(1-z)^k$.

19.3 A compact verification script

The following pseudocode is not a proof, but it is useful for regression tests when translating the atlas into Lean, a CAS, or a library implementation.

```

for n from 0 to N:
  s := popcount(n)
  eps := 1 - 2*(s mod 2)

  assert v2(factorial(n)) == n - s
  assert v2(binomial(2*n,n)) == s

  oddPascal := count k in [0,n] with binomial(n,k) odd
  assert oddPascal == 2^s
  assert (oddPascal mod 3) - 1 == (s mod 2)

  sparseParity := sum binomial(n,2^j) over all 2^j <= n
  assert (-1)^sparseParity == eps

for m from 0 to M:
  P := sum_{n=0}^{2^m-1} eps(n)*z^n
  assert P == product_{j=0}^{m-1} (1-z^(2^j))

  for r from 0 to m-1:
    assert sum_{n=0}^{2^m-1} eps(n)*n^r == 0

  assert sum_{n=0}^{2^m-1} eps(n)*n^m
    == (-1)^m * factorial(m) * 2^(m*(m-1)/2)

```

19.4 A dependency-aware formalization order

A conservative Lean implementation should formalize the finite algebraic spine before introducing analysis.

1. **Binary layer.** Prove recurrence, block concatenation, XOR, reflection, and the two-kernel from a chosen binary-weight primitive.
2. **Valuation layer.** Establish the base-two Legendre formula, central-binomial valuation, Catalan valuation, and Kummer carry identity.

3. **Pascal/Frobenius layer.** Formalize Lucas support, Glaisher count, sparse probes, the modulo-three consequence, multinomial support, and Boolean Möbius inversion.
4. **Polynomial layer.** Define P_m , prove the finite product, reciprocity, cyclotomic multiplicities, and signed Pascal-support products over a suitable coefficient ring.
5. **Difference and moment layer.** First isolate finite weighted polynomial functionals and their coefficient-extraction rules over commutative semirings. Then derive the ring-valued mixed differences, Prouhet cancellation, exact top moments, composition formulas, and sparse moments. No topology or characteristic-zero hypothesis belongs in this algebraic layer unless a surviving scalar is explicitly cancelled.
6. **Formal-series layer.** Add the Mahler equations, Euler transform, ruler convolution, Bell/determinant forms, and the Artin–Schreier equation over $\mathbb{F}_2[[z]]$.
7. **Analytic layer.** Only then introduce compact-uniform convergence, Fourier evaluation, Riesz measures, Mellin transforms, and infinite products.

Possible module boundaries are listed below. Entries marked *Delivered* name modules that already appear in the inventory of [Section 19.1](#); they are repeated here to record which proposed boundary each one realizes, and are not additional modules.

Proposed module	Main declarations
ThueMorseValuation.lean, CentralBinomialValuation.lean	Factorial, Catalan, successor, and carry formulae, together with the exhaustive three-theorem central-binomial valuation/sign bridge. <i>Delivered</i> .
ThueMorsePascal.lean	Lucas submasks, Glaisher count, sparse probes, modulo-three formula. <i>Delivered</i> (in Parity.lean, ThueMorseBinomialLog.lean, ThueMorseLucasSupport.lean).
ThueMorseMultinomial.lean	Frobenius factorization and $O_r(n) = r^{\text{wt}_2(n)}$. <i>Delivered</i> (via Kummer’s criterion rather than Frobenius).
ThueMorseBooleanMobius.lean	Submask incidence algebra and inverse parity matrix. <i>Delivered</i> (poset form; the matrix form is its row-by-row reading).
ThueMorseGeneratingProduct.lean	P_m , cyclotomic multiplicities, Mahler equations. <i>Delivered</i> (ThueMorseBlockAlgebra.lean, ThueMorseMahler.lean).
ThueMorseProuhetMoments.lean	Difference operators, all moments, centered parity. <i>Delivered</i> (ThueMorseMoments.lean, ThueMorseMixedDifference.lean, ThueMorseSymmetricDifference.lean).
ThueMorseSparseProuhet.lean	Arbitrary Pascal-support products and sparse moments. <i>Delivered</i> (plus ThueMorseSparseMoments.lean).
ThueMorseFormalSeriesF2.lean	Artin–Schreier equation and explicit solution. <i>Delivered</i> (ThueMorseArtinSchreier.lean, with uniqueness and the integer lift).
ThueMorseFourier.lean	Finite DFT, Walsh transform, product-to-sum identities. <i>Delivered</i> across ThueMorseComplexExponential.lean, ThueMorseComplexProductBridge.lean, ThueMorseFourier.lean, ThueMorseSineProduct.lean, LacunaryProductToSum.lean, LacunaryProductToSumCosine.lean, and ThueMorseFourierInversion.lean: the raw weighted and signed exponential transforms for every complex parameter and affine translate, the total complex half-angle and sinc normalizations, the Walsh and Riesz factorizations, the affine real sine and cosine forms, both real lacunary product-to-sum identities, and Fourier inversion are formal. Only direct arbitrary-complex declarations in the two rescaled product-to-sum forms remain unbundled.
ThueMorseDirichlet.lean	Boundary flatness, Mellin integral, entire continuation. <i>Delivered</i> (ThueMorseBoundaryFlatness.lean, ThueMorseDirichlet.lean, ThueMorseMellin.lean, ThueMorseEntireContinuation.lean).
ThueMorseRarefied.lean	Root-of-unity filters, modulo-three closed forms, odd-modulus matrix recursion. <i>Delivered in full</i> (ThueMorseArithmeticFilter.lean, ThueMorseRarefied.lean).
ThueMorseWords.lean	Uniform recurrence, aperiodicity, and the elementary complexity bounds. <i>Delivered in full</i> (ThueMorseAperiodicity.lean, ThueMorseOverlapFree.lean, ThueMorseRecurrence.lean, ThueMorseComplexity.lean).

Every planned layer has now been delivered, in most cases in greater generality than proposed.

19.5 Proof-dependency graph

The central dependencies can be summarized as

$$\begin{aligned}
 \text{binary digits} &\implies \begin{cases} \text{recursion and XOR,} \\ \text{Lucas/Frobenius support,} \\ \text{finite product } P_m, \end{cases} \\
 P_m &\implies \begin{cases} \text{Mahler and Euler equations,} \\ \text{Fourier/Walsh products,} \\ \text{order-}m \text{ zero at 1,} \end{cases} \\
 \text{order-}m \text{ zero} &\implies \begin{cases} \text{Prouhet cancellation,} \\ \text{all finite moments,} \\ \text{iterated-prefix flatness,} \\ \text{Riesz-product normalization.} \end{cases}
 \end{aligned} \tag{19.1}$$

Legendre/Kummer form a parallel arithmetic branch; Lucas and the signed support product join that branch to the finite-difference branch at sparse Prouhet identities. Boundary flatness of the infinite product then opens the Mellin/Dirichlet branch.

20 Worked examples and consistency checks

20.1 One index through many representations

Take $n = 13 = (1101)_2$. The occupied positions are 0, 2, 3, so

$$\text{wt}_2(13) = 3, \quad \tau(13) = 1, \quad \varepsilon_{13} = -1. \tag{20.1}$$

The floor, factorial, and central-binomial routes give

$$\sum_{j \geq 1} \left\lfloor \frac{13}{2^j} \right\rfloor = 6 + 3 + 1 = 10, \tag{20.2}$$

$$(-1)^{13 - \nu_2(13!)} = (-1)^3 = -1, \tag{20.3}$$

$$\nu_2 \binom{26}{13} = 3. \tag{20.4}$$

The distinguished Pascal entries are

$$\binom{13}{1} = 13, \quad \binom{13}{2} = 78, \quad \binom{13}{4} = 715, \quad \binom{13}{8} = 1287. \tag{20.5}$$

Their parities are 1, 0, 1, 1, exactly the four low bits of 13, and their sum 2093 is odd. Row 13 has

$$\nu(13) = 2^3 = 8 \tag{20.6}$$

odd entries, so $\nu(13) \bmod 3 = 2$ and the logarithm-free formula gives $\tau(13) = 2 - 1 = 1$. In three variables, exactly

$$O_3(13) = 3^3 = 27 \tag{20.7}$$

coefficients of $(x + y + z)^{13}$ are odd.

The submasks are

$$0, 1, 4, 5, 8, 9, 12, 13, \tag{20.8}$$

and their signed support polynomial is

$$R_{13}(z) = (1 - z)(1 - z^4)(1 - z^8). \quad (20.9)$$

Its derivatives of orders 0, 1, 2 at one vanish, and

$$\sum_{k \prec_2 13} \varepsilon_k k^3 = -3! 2^{0+2+3} = -192. \quad (20.10)$$

All these evaluations recover the same three occupied bits through different algebraic interfaces.

20.2 An all-order dyadic moment

For $m = 3$ and $r = 5$, direct summation gives

$$M_{3,5} = \sum_{n=0}^7 \varepsilon_n n^5 = -6720. \quad (20.11)$$

The positive compositions of 5 into three parts are

$$(1, 1, 3), (1, 2, 2), (1, 3, 1), (2, 1, 2), (2, 2, 1), (3, 1, 1).$$

Formula (9.27) gives

$$M_{3,5} = -5! \sum_{\substack{p_0+p_1+p_2=5 \\ p_j \geq 1}} \frac{2^{p_1+2p_2}}{p_0! p_1! p_2!} = -6720. \quad (20.12)$$

No recursion through lower moments is needed.

20.3 A Fourier sample

For $m = 4$, every even DFT frequency vanishes. At $k = 1$,

$$\text{DFT}_{2^4}[\varepsilon][1] = (2i)^4 e^{-15\pi i/16} \sin \frac{\pi}{2} \sin \frac{\pi}{4} \sin \frac{\pi}{8} \sin \frac{\pi}{16}. \quad (20.13)$$

The factorization exposes both the exact phase and the hierarchy of four dyadic scales, which direct summation of sixteen complex terms obscures.

20.4 Autocorrelation from the recursion

Starting with $\eta(0) = 1$, the fixed-point equation at one gives $\eta(1) = -1/3$. Then

$$\begin{aligned} \eta(2) &= \eta(1) = -\frac{1}{3}, \\ \eta(3) &= -\frac{\eta(1) + \eta(2)}{2} = \frac{1}{3}, \\ \eta(5) &= -\frac{\eta(2) + \eta(3)}{2} = 0. \end{aligned}$$

These are simultaneously combinatorial averages and Fourier coefficients of the singular continuous Thue–Morse measure.

20.5 A complete four-bit block

At $m = 4$,

$$\begin{aligned} P_4(z) &= 1 - z - z^2 + z^3 - z^4 + z^5 + z^6 - z^7 \\ &\quad - z^8 + z^9 + z^{10} - z^{11} + z^{12} - z^{13} - z^{14} + z^{15} \\ &= (1 - z)(1 - z^2)(1 - z^4)(1 - z^8). \end{aligned} \quad (20.14)$$

The factor $(1 - z)^4$ divides the product, so

$$\sum_{n < 16} \varepsilon_n n^r = 0 \quad (0 \leq r < 4), \quad \sum_{n < 16} \varepsilon_n n^4 = 4! 2^{\binom{4}{2}} = 1536, \quad (20.15)$$

the sign being positive because $(-1)^4 = 1$. The residue-class sums modulo three are, by [theorem 15.1](#) at level $m = 4$,

$$(A_{4,0}^{(3)}, A_{4,1}^{(3)}, A_{4,2}^{(3)}) = (6, -3, -3). \quad (20.16)$$

Their total is zero, as block balance requires, yet the distribution among residue classes is already strongly skewed toward the multiples of three.

20.6 A carry-corrected addition

Take $a = 13 = (1101)_2$ and $b = 7 = (0111)_2$. Then $\varepsilon_{13} = \varepsilon_7 = -1$, while $a + b = 20 = (10100)_2$ has $\varepsilon_{20} = +1$. Kummer’s formula [\(4.9\)](#) counts

$$\nu_2 \binom{20}{13} = \text{wt}_2(13) + \text{wt}_2(7) - \text{wt}_2(20) = 3 + 3 - 2 = 4 \quad (20.17)$$

carries, and the correction $(-1)^4 = 1$ in [\(4.10\)](#) restores multiplicativity: $\varepsilon_{20} = (-1)(-1)(+1)$. The example shows why a yes/no record of “was there a carry” cannot suffice: the sign needs the carry count with propagation, taken modulo two.

21 Concluding synthesis

The formula in Section 11.8 of the Fabius–Rvachev exposition is not an isolated trick. It lies at the intersection of four structures:

binary digits $\longleftrightarrow$ Frobenius/Pascal support $\longleftrightarrow$ Boolean-lattice characters $\longleftrightarrow$ dyadic finite differences
--

(21.1)

The signed Thue–Morse sequence is the common coefficient system of this translation dictionary.

The arithmetic formulas hide digit parity in factorial valuations, central binomial coefficients, Catalan numbers, and carry counts. The Pascal formulas recover it either globally, from the number of odd row entries, or sparsely, from the power-of-two columns. The modulo-three identity shows that the exact base-two logarithm is optional: the exponent parity can also be read from the cycle of powers of two in a finite residue field. Multinomial support turns the same mechanism into a whole $r^{\text{wt}_2(n)}$ family.

The product $P_m(z) = \prod_{j < m} (1 - z^{2^j})$ is the central unifying engine. At $z = 1$ it has an exact high-order zero and produces Prouhet cancellation. At $z = e^t$ it produces every finite moment. At roots of unity it produces the dyadic DFT. On the Boolean cube it is a pure Walsh character. After taking absolute squares it produces Riesz densities. After division by $(1 - z)^k$ it produces iterated-prefix grids and their terminal flatness. After the limit $m \rightarrow \infty$, it becomes a Mahler function, an Euler transform of the ruler sequence, and the Mellin kernel of an entire Dirichlet family.

The added Mellin theory sharpens the analytic picture: the infinite product is not merely small near one, but flat beyond every algebraic order, with a quadratic leading term in $\log(1/t)$. That single estimate explains why all potential Mellin poles disappear and why the shifted Thue–Morse Dirichlet series has zeros at every nonpositive integer. Its derivative at zero then recovers the logarithms of the gamma-like products generated by Thue–Morse exponents.

For the Fabius project, the most useful lesson is proof-interface diversity. No one formula dominates every task. Block multiplicativity is ideal for dyadic reindexing; the carry cocycle for ordinary sums; Pascal support for characteristic-two arguments; finite differences for cancellation; Bell polynomials for all moments; Walsh characters for finite harmonic analysis; and the Mellin product for analytic continuation. Formalizing several of these interfaces would not duplicate mathematics: it would let later Fabius proofs choose the representation adapted to the operation being performed.

A The first 64 terms

The rows display $\tau(n)$; replacing 0 by +1 and 1 by −1 gives ε_n .

Range	successive values							
0–7	0	1	1	0	1	0	0	1
8–15	1	0	0	1	0	1	1	0
16–23	1	0	0	1	0	1	1	0
24–31	0	1	1	0	1	0	0	1
32–39	1	0	0	1	0	1	1	0
40–47	0	1	1	0	1	0	0	1
48–55	0	1	1	0	1	0	0	1
56–63	1	0	0	1	0	1	1	0

Every row of length eight is a copy or complement according to the sign of its block index, as predicted by (2.6).

B A compact identity dictionary

For quick reference, the principal exact translations are

$$\begin{aligned}
 \tau(n) &\longleftrightarrow \varepsilon_n = 1 - 2\tau(n), \\
 \varepsilon_n &\longleftrightarrow \varepsilon_n, \\
 \text{wt}_2(n) &\longleftrightarrow n - \nu_2(n!), \\
 \text{wt}_2(n) &\longleftrightarrow \nu_2\binom{2n}{n}, \\
 b_j(n) &\longleftrightarrow \binom{n}{2^j} \bmod 2, \\
 2^{\text{wt}_2(n)} &\longleftrightarrow \#\{k : \binom{n}{k} \text{ odd}\}, \\
 \varepsilon_n &\longleftrightarrow [z^n] \prod_{j \geq 0} (1 - z^{2^j}), \\
 \sum_{n < 2^m} \varepsilon_n f(x + n) &\longleftrightarrow (-1)^m \Delta_1 \Delta_2 \cdots \Delta_{2^{m-1}} f(x),
 \end{aligned}$$

$$P_m(e^{-2\pi i k/2^m}) \longleftrightarrow \text{DFT}_{2^m}[\varepsilon][k],$$

$$E(e^{-t}) \longleftrightarrow \Gamma(s)D(s, a) \text{ under Mellin transformation.}$$

Every arrow denotes an exact identity or invertible conversion, not an asymptotic analogy.

C Draft-to-section concordance

The four source drafts overlap substantially. The following table records the main distinctive material retained from each, together with additions made in the unified treatment.

Source	Distinctive contributions represented in this atlas
Draft 1	Carry and XOR viewpoints; Boolean Möbius inversion and inverse Pascal parity matrix; ruler Euler transform; partition, Bell, and Hessenberg determinant formulae; finite-difference positivity and analytic transforms.
Draft 2	Sparse power-of-two Pascal columns; all-order dyadic moments; automaton/tensor/Walsh forms; exact DFT and Riesz products; master infinite products; base- q root-of-unity extension.
Draft 3	Multinomial support and exact r -nomial logarithms; sparse Prouhet theorem on arbitrary Pascal rows; autocorrelation recurrences; Artin–Schreier series over $\mathbb{F}_2$; Catalan and fractional-part formulae.
Draft 4	Logarithm-free modulo-three Pascal formula; sinc-product interval signs; evil/odious enumerators; integer algebraic lift; trigonometric product-to-sum identities; explicit iterated-prefix/Fabius bridge.
Unified additions	Dyadic weight-distribution chapter; centered-moment parity selection; all sparse Pascal-support moments; elementary boundary asymptotic for $E(e^{-t})$; entire shifted Dirichlet family, its nonpositive-integer zeros, and its derivative bridge to exponent products; sharp first moment for the general base- q character.
Parallel atlas 1	Self-contained natural-boundary proof; supertile arguments for uniform recurrence and the linear complexity bound $p(\ell) < 8\ell$; the literature panorama of arithmetic subsequences (primes, squares, cubes, level of distribution).
Parallel atlas 2	Gelfond’s sharp uniform exponent with the lower bound at cube roots of unity; the odd-modulus matrix recursion; the autocorrelation Mahler equation; subtraction/borrow cocycle; weighted digit enumerator; Hankel/Padé framing of irrationality measure.
Parallel atlas 3	Exact modulo-three rarefaction vectors with the two-scale constants; Newman positivity framing; complement/reversal closure of the factor language; critical exponent statement; Mahler value transcendence and $\mu = 2$ exposition; four-bit block and carry-corrected worked examples.
Second-pass additions	Even-modulus collapse of rarefied sums from exact prefix identities; the characteristic-zero/characteristic-two reading of the lacunary digit-sum series (6.20); Morse–Hedlund lower complexity bound paired with the elementary upper bound.

References

- [1] V. Reshetnikov, *ProveIt: Analysis/FabiusFunction*, GitHub repository, source snapshot 8e7aa582661f2828302f70748efe20d02d8ae9c7, 2026.
- [2] E. Prouhet, Mémoire sur quelques relations entre les puissances des nombres, *Comptes rendus de l'Académie des Sciences de Paris* **33** (1851), 225.
- [3] A. Thue, *Über unendliche Zeichenreihen*, Norske Videnskabers Selskabs Skrifter, I. Matematisk-naturvidenskabelig Klasse, no. 7 (1906), 1–22.
- [4] M. Morse, Recurrent geodesics on a surface of negative curvature, *Transactions of the American Mathematical Society* **22** (1921), 84–100.
- [5] E. E. Kummer, Über die Ergänzungssätze zu den allgemeinen Reciprocitätsgesetzen, *Journal für die reine und angewandte Mathematik* **44** (1852), 93–146.
- [6] N. J. Fine, Binomial coefficients modulo a prime, *American Mathematical Monthly* **54** (1947), 589–592.
- [7] R. Meštrović, Lucas’ theorem: its generalizations, extensions and applications (1878–2014), arXiv:1409.3820, 2014.
- [8] K. Mahler, Arithmetische Eigenschaften der Lösungen einer Klasse von Funktionalgleichungen, *Mathematische Annalen* **101** (1929), 342–366.
- [9] D. R. Woods, Elementary problem E2692, *American Mathematical Monthly* **85** (1978), 48.
- [10] D. Robbins, Solution to elementary problem E2692, *American Mathematical Monthly* **86** (1979), 394–395.
- [11] J.-P. Allouche and J. Shallit, The ubiquitous Prouhet–Thue–Morse sequence, in C. Ding, T. Helleseth, and H. Niederreiter (eds.), *Sequences and Their Applications*, Springer, London, 1999, 1–16.
- [12] J.-P. Allouche and J. Shallit, *Automatic Sequences: Theory, Applications, Generalizations*, Cambridge University Press, Cambridge, 2003.
- [13] G. Christol, T. Kamae, M. Mendès France, and G. Rauzy, Suites algébriques, automates et substitutions, *Bulletin de la Société Mathématique de France* **108** (1980), 401–419.
- [14] M. Queffélec, *Substitution Dynamical Systems—Spectral Analysis*, 2nd ed., Lecture Notes in Mathematics 1294, Springer, Berlin, 2010.
- [15] M. Baake and U. Grimm, The singular continuous diffraction measure of the Thue–Morse chain, *Journal of Physics A: Mathematical and Theoretical* **41** (2008), 422001.
- [16] J.-P. Allouche, S. Riasat, and J. Shallit, More infinite products: Thue–Morse and the gamma function, *Ramanujan Journal* **49** (2019), 115–128; arXiv:1709.03398.
- [17] P. Borwein and C. Ingalls, The Prouhet–Tarry–Escott problem revisited, *L’Enseignement Mathématique* **40** (1994), 3–27.
- [18] Wikipedia, Prouhet–Tarry–Escott problem, https://en.wikipedia.org/wiki/Prouhet%E2%80%93Tarry%E2%80%93Escott_problem, accessed 26 August 2026.
- [19] M. Morse and G. A. Hedlund, Symbolic dynamics, *American Journal of Mathematics* **60** (1938), 815–866.

- [20] F. Carlson, Über Potenzreihen mit ganzzahligen Koeffizienten, *Mathematische Zeitschrift* **9** (1921), 1–13.
- [21] J. P. Bell, M. Coons, and E. Rowland, The rational–transcendental dichotomy of Mahler functions, *Journal of Integer Sequences* **16** (2013), Article 13.2.10.
- [22] J.-P. Allouche, J. Peyrière, Z.-X. Wen, and Z.-Y. Wen, Hankel determinants of the Thue–Morse sequence, *Annales de l’Institut Fourier* **48** (1998), 1–27.
- [23] G.-N. Han, Hankel determinant calculus for the Thue–Morse and related sequences, arXiv:1406.1586, 2014.
- [24] Y. Bugeaud and G.-N. Han, A combinatorial proof of the non-vanishing of Hankel determinants of the Thue–Morse sequence, *Electronic Journal of Combinatorics* **21**(3) (2014), Paper 3.26.
- [25] Y. Bugeaud, On the rational approximation to the Thue–Morse–Mahler numbers, *Annales de l’Institut Fourier* **61** (2011), 2065–2076.
- [26] Y. Bugeaud and M. Queffélec, On the rational approximation to the binary Thue–Morse–Mahler number, *Journal of Integer Sequences* **16** (2013), Article 13.2.1.
- [27] Y. Bugeaud, G.-N. Han, Z.-Y. Wen, and J.-Y. Yao, Hankel determinants, Padé approximations, and irrationality exponents, arXiv:1503.02797, 2015.
- [28] A. O. Gelfond, Sur les nombres qui ont des propriétés additives et multiplicatives données, *Acta Arithmetica* **13** (1968), 259–265.
- [29] M. Queffélec, Questions around the Thue–Morse sequence, *Uniform Distribution Theory* **13** (2018), 1–25.
- [30] D. J. Newman, On the number of binary digits in a multiple of three, *Proceedings of the American Mathematical Society* **21** (1969), 719–721.
- [31] J. Coquet, A summation formula related to the binary digits, *Inventiones Mathematicae* **73** (1983), 107–115.
- [32] M. Drmota and M. Skalba, Rarified sums of the Thue–Morse sequence, *Transactions of the American Mathematical Society* **352** (2000), 609–642.
- [33] M. Drmota and T. Stoll, Newman’s phenomenon for generalized Thue–Morse sequences, *Discrete Mathematics* **308** (2008), 1191–1208.
- [34] J. Shallit, Rarefied Thue–Morse sums via automata theory and logic, arXiv:2302.09436, 2023.
- [35] C. Mauduit and J. Rivat, Sur un problème de Gelfond: la somme des chiffres des nombres premiers, *Annals of Mathematics* **171** (2010), 1591–1646.
- [36] M. Drmota, C. Mauduit, and J. Rivat, Normality along squares, *Journal of the European Mathematical Society* **21** (2019), 507–548.
- [37] L. Spiegelhofer, The level of distribution of the Thue–Morse sequence, *Compositio Mathematica* **156** (2020), 2560–2587.
- [38] L. Spiegelhofer, Thue–Morse along the sequence of cubes, arXiv:2308.09498, 2023.

Part II

Finite-Block Calculus for the Fabius–Rvachev–Thue–Morse System

Abstract

The documentation tree under `Analysis/FabiusFunction` already contains an unusually broad and internally cross-checked theory of the Thue–Morse sequence, the bounded Fabius distribution function, Rvachev’s compactly supported up-function, dyadic q -binomial arithmetic, infinite sinc products, and the lower-Lambert endpoint phase. This report audits the living manuscripts, the imported primary literature, the active Fourier-decay drafts, and the frozen historical variants, then develops a new finite-block calculus that joins several strands which had previously remained adjacent rather than identical.

The central observation is exact. A single length- 2^m signed Thue–Morse block polynomial, under an oscillatory substitution, is a phase-and-scale normalization of the first m factors of the Fourier sinc product; under a decaying substitution, the same polynomial is a normalization of the first m factors of the negative-Laplace product used in endpoint asymptotics. The omitted tail is not merely small: it is a scaled copy of the full product. This yields (i) a convergent all-orders expansion whose coefficients are zeta values arranged as a Lambert series with base $q = 1/4$; (ii) explicit, sharp-radius remainder bounds; (iii) arbitrary-order finite Thue–Morse approximants to the Fourier transform; (iv) a q -binomial Richardson scheme whose order- p extrapolant converges like $4^{-(p+1)m}$; (v) an exact random-tail decomposition and a weak differential expansion of the up-density in terms of finite uniform splines; (vi) a Lambert- W law for the number of product factors needed on the canonical endpoint saddle scale; and (vii) a Barnes-like tower built from derivatives of the entire shifted Thue–Morse Dirichlet function at its trivial zeros.

All stated theorems are proved here. High-precision computations independently check the finite identities, coefficients, asymptotic orders, and truncation laws. The phrase “new result” means new to the inspected repository corpus and derived in this report; no claim of global publication priority is made without a dedicated literature review beyond the searches recorded here.

Status discipline. The report uses three levels. Statements explicitly attributed to the repository or to the literature are background. Statements labeled “Result derived in this report” are supplied with proofs. Numerical experiments are validation, not substitutes for proof. Several final research directions are labeled as conjectures or formalization obligations.

Formalization crosswalk. “Result derived in this report” is a statement about the inspected *prose* corpus and says nothing about the Lean library. In fact several of the results below already have formal counterparts in `Analysis/FabiusFunction/Lean/FabiusFunction/`, in some cases proved there in greater generality than they are stated here. Wherever this is so, a paragraph headed *Formal version* follows the statement and gives the exact declaration names, typeset as `Fabius.thueMorseSign`, together with the module they live in; where the Lean statement is weaker (for example coefficientwise rather than analytically evaluated) or stronger (arbitrary ring, arbitrary base, arbitrary interpolation ratio) than the statement here, the paragraph says which. Absence of such a paragraph means only that no formal counterpart was located, and must not be read as a claim that none exists. A paper proof is in any case not evidence of formalization, and a citation to a more general Lean theorem is not a claim that the exact statement printed here has been formalized.

1 Corpus audit and research positioning

1.1 What was inspected

The repository’s documentation is not one article but a layered research environment. The living primary exposition separates the bounded cumulative distribution F , the even bump up, and the signed global extension, and develops their construction, probability model, Fourier products, dyadic arithmetic, and corrected endpoint asymptotics. The consolidated frontier volume records non-formalized deductions and gaps; the Thue–Morse Atlas develops a large formula network; the active Fourier-decay directory contains twelve mutually auditing manuscripts; the non-elementarity article and Lean walkthrough provide logical and formal context; and the papers directory contains source TeX for relevant primary literature. The archive is explicitly a provenance-preserving collection of frozen versions, with byte-identical or otherwise redundant entries pruned from the historical catalogue [1, 2, 3, 4, 5].

The audit protocol was therefore:

1. read the current living manuscripts as the authoritative mathematical state;
2. compare the active Fourier-decay drafts because they contain genuinely different hypotheses, corrections, and audit conclusions;
3. inspect each imported paper represented by TeX source;
4. use the archive’s commit-level provenance map to deduplicate superseded historical copies, while retaining every uniquely described line of development in the research map;
5. search the corpus for the candidate identities and coefficient patterns before calling them new to the corpus.

A complete source ledger appears in [section C](#).

1.2 Established pillars

The following table compresses the established state which this report uses but does not re-prove in full.

Pillar	Established content	Remaining junction exploited here
Thue–Morse	$\varepsilon_n = (-1)^{\text{wt}_2(n)}$, block products, Prouhet cancellation, Mahler equations, Walsh and lacunary products, shifted entire Dirichlet theory	Turn one finite block into a certified approximation engine for both Fourier and Laplace products
Fabius/up probability	up as a compactly supported probability density and F as its rescaled CDF; finite atomic/uniform convolution models	Identify the omitted convolution tail as an exact scaled copy and expand it to all orders
Fourier analysis	$\mathcal{F}_{2\pi}[\text{up}]$ as a dyadic sinc product; shell factorization; pointwise, metric, averaged, and extremal decay regimes	Analyze truncation in the small tail variable, which is complementary to large-frequency shell dynamics
Dyadic arithmetic	exact values at dyadic rationals, q -Pochhammer and Gaussian-binomial formulae at base $1/2$, Thue–Morse coefficient extraction	Produce a new Gaussian-binomial role at base $q = 1/4$: geometric-grid extrapolation weights
Endpoint asymptotics	negative-Laplace product, periodic lower-Lambert phase, all-orders saddle expansion, $\lambda = -W_{-1}(-x \log 2)/\log 2$	Quantify the finite-product depth on the canonical Lambert scale and express it through one finite Thue–Morse block
Formalization	extensive Lean construction and exact proof correspondence; explicit quarantine of unformalized claims	Isolate short algebraic and analytic lemmas suitable for promotion

1.3 The missing commutative bridge

The corpus contains all four vertices of the following conceptual square:

$$\begin{array}{ccc}
 \text{finite Thue–Morse blocks} & \longrightarrow & \text{finite dyadic products} \\
 \downarrow & & \downarrow \\
 \text{exact dyadic arithmetic} & \longrightarrow & \text{Fourier/Laplace asymptotics.}
 \end{array}$$

What was not isolated as one theorem package is that the top horizontal arrow has two exact specializations, oscillatory and decaying, and that the right vertical arrow has a self-similar tail with an analytic expansion in the geometric parameter 4^{-m} . Once this is written down, zeta values, Lambert series, Gaussian binomials, random-tail moments, and the Lambert- W endpoint phase become parts of one finite-block calculus rather than separate phenomena.

1.4 Priority and uncertainty

Keyword searches of the inspected TeX corpus did not locate the all-orders 4^{-m} tail formula, the sharp-radius remainder bound, or the q -binomial extrapolation scheme in the form proved below. Searches of arXiv and publisher records located the classical Fabius, Thue–Morse, lacunary-product, automatic-sequence, arithmetic, Dirichlet-series, and Lambert- W sources used for context [6, 7, 8, 9, 10, 11], but did not constitute an exhaustive priority search. Accordingly, every novelty statement in this report is corpus-relative.

2 Objects, normalizations, and exact background

2.1 The signed block polynomial

Let

$$\varepsilon_n = (-1)^{\text{wt}_2(n)}, \quad P_m(z) = \sum_{n=0}^{2^m-1} \varepsilon_n z^n = \prod_{j=0}^{m-1} (1 - z^{2^j}), \quad (2.1)$$

where $\text{wt}_2(n)$ is the binary digit sum. The product identity is the finite form of the Thue–Morse generating-product law. Its zero of order m at $z = 1$ is exactly the Prouhet cancellation of moments of degrees below m .

Formal version. The product identity in (2.1) is `Fabius.prod_one_sub_pow_eq_sum_thueMorseSign (ThueMorseBooleanCube.lean)`, which is more general than the form used here: it is stated for z in an arbitrary commutative ring rather than being restricted to complex z , and is proved once by a Boolean-cube expansion. The complex substitution step is now also exposed directly by four theorems in `ThueMorseComplexExponential.lean`: `Fabius.sum_range_two_pow_binaryWeight_cexp` and its affine companion `Fabius.sum_range_two_pow_binaryWeight_cexp_affine`, together with the signed forms `Fabius.sum_range_two_pow_thueMorseSign_cexp` and `Fabius.sum_range_two_pow_thueMorseSign_cexp_affine`. Thus the raw finite exponential block and an arbitrary common complex exponent shift are formal for every complex parameter, including $m = 0$. The substitutions below — exponent parameter $it/2^{m-1}$ in [theorem 3.1](#) and $-s/2^m$ in [theorem 3.3](#) — are instances of the signed theorem. This module is the raw finite layer; `ThueMorseComplexProductBridge.lean` separately supplies the total complex sinc and negative-Laplace normalizations, including their removable origins. Neither module asserts an infinite product.

2.2 Fourier and Laplace products

Use $\text{sinc}_{\text{rad}} z = \sin z / z$ with $\text{sinc}_{\text{rad}} 0 = 1$, and define

$$Q_m(t) = \prod_{k=1}^m \text{sinc}_{\text{rad}}\left(\frac{t}{2^k}\right), \quad Q_\infty(t) = \prod_{k=1}^{\infty} \text{sinc}_{\text{rad}}\left(\frac{t}{2^k}\right), \quad (2.2)$$

$$G_m(-s) = \prod_{k=1}^m \frac{1 - e^{-s/2^k}}{s/2^k}, \quad G(-s) = \prod_{k=1}^{\infty} \frac{1 - e^{-s/2^k}}{s/2^k}. \quad (2.3)$$

The products converge locally uniformly and define entire functions after removable values at zero are supplied.

With the Fourier convention

$$\widehat{h}_{2\pi}(\xi) = \int_{\mathbb{R}} h(x) e^{-2\pi i x \xi} dx,$$

the normalized Rvachev density satisfies

$$\widehat{\text{up}}_{2\pi}(\xi) = Q_\infty(2\pi\xi) = \prod_{j=0}^{\infty} \text{sinc}_{\text{rad}}\left(\frac{\pi\xi}{2^j}\right). \quad (2.4)$$

The bounded Fabius distribution is related by

$$F'(x) = 2 \text{up}(2x - 1), \quad F(x) = \int_{-1}^{2x-1} \text{up}(y) dy \quad (0 \leq x \leq 1). \quad (2.5)$$

These conventions agree with the repository’s primary exposition [1, 6].

Formal versions, and the convention check. Both normalizations used above are the ones carried by the Lean library, which matters because the corpus contains more than one Fourier convention and a crosswalk across a mismatch would be worthless. The Lean transform `Fabius.rvachevFourier` is defined as $\int_{\mathbb{R}} \text{up}(t) e^{-2\pi i t z} dt$ — the convention of (2.4) — and `Fabius.rvachevFourier_eq_product (FourierProduct.lean)` proves it equal to `Fabius.rvachevFourierProduct`, namely $\prod_{n \geq 0} \text{sinc}_{\text{rad}}(\pi z / 2^n)$. That is (2.4) verbatim. Likewise G in (2.3) is the Lean object `Fabius.complexGeneratingFunction`: that the entire half-moment generating function evaluated at $-s$ equals the dyadic product $\prod_{k \geq 1} (1 - e^{-s/2^k}) / (s/2^k)$ of (2.3), with the same removable value at each vanishing factor, is `Fabius.complexGeneratingFunction_neg_eq_tprod (NegativeLaplaceVertical.lean)`, for every $s \in \mathbb{C}$.

2.3 The lower-Lambert phase

Write $L = \log 2$. For small $x > 0$, the canonical endpoint phase variable is

$$\lambda(x) = -\frac{W_{-1}(-Lx)}{L}, \quad x = \lambda(x) 2^{-\lambda(x)}. \quad (2.6)$$

The associated leading saddle scale is

$$s_0(x) = 2^{\lambda(x)} = \frac{\lambda(x)}{x}. \quad (2.7)$$

The repository’s endpoint theory adds periodic corrections in $\log_2 s$ and an all-orders saddle expansion. Our use of (2.7) below is computational: it measures how many factors of the exact product are needed near that scale. It is not a replacement for the periodically corrected saddle analysis.

3 One Thue–Morse block, two exact product bridges

3.1 Oscillatory specialization: the Fourier bridge

Theorem 3.1 (Exact finite Thue–Morse–sinc identity). *For every integer $m \geq 1$ and every $t \in \mathbb{C}$,*

$$P_m\left(e^{it/2^{m-1}}\right) = (-i)^m t^m 2^{-\binom{m}{2}} e^{it(1-2^{-m})} Q_m(t). \quad (3.1)$$

Equivalently, with the removable value at $t = 0$ understood,

$$Q_m(t) = i^m 2^{\binom{m}{2}} t^{-m} e^{-it(1-2^{-m})} \sum_{n=0}^{2^m-1} \varepsilon_n e^{int/2^{m-1}}. \quad (3.2)$$

Result derived in this report.

Proof. Insert $z = e^{it/2^{m-1}}$ in the product form of P_m and reverse the index:

$$P_m(z) = \prod_{k=0}^{m-1} \left(1 - e^{it/2^k}\right).$$

For each factor use

$$1 - e^{ix} = (-i)xe^{ix/2} \operatorname{sinc}_{\text{rad}}(x/2).$$

The product of the linear terms is

$$(-i)^m t^m 2^{-\sum_{k=0}^{m-1} k} = (-i)^m t^m 2^{-\binom{m}{2}}.$$

The product of the phases is

$$\exp\left(\frac{it}{2} \sum_{k=0}^{m-1} 2^{-k}\right) = e^{it(1-2^{-m})},$$

and the sinc factors are $\operatorname{sinc}_{\text{rad}}(t/2), \dots, \operatorname{sinc}_{\text{rad}}(t/2^m)$. This gives (3.1); the sum form follows from (2.1). $\square$

Formal version. The total denominator-free form of (3.1) is `Fabius.thueMorseBlock_cexp_eq_sincPrefix_of_pos` (`ThueMorseComplexProductBridge.lean`); the stronger all-level scaling is `Fabius.thueMorseBlock_cexp_eq_sincPrefix`. The Lean prefix `Fabius.shiftedComplexSincPrefix` is exactly $Q_m(t)$. The primary equality is over $\mathbb{C}$ and includes $t = 0$, while `Fabius.shiftedComplexSincPrefix_apply_zero` gives $Q_m(0) = 1$. For $t \neq 0$, the quotient normalization (3.2) is `Fabius.shiftedComplexSincPrefix_eq_thueMorseBlock_cexp_of_pos`. The supporting complex-argument sine product and half-angle identity are `Fabius.sum_thueMorseSign_cexp_eq_sin_prod` and `Fabius.one_sub_cexp_mul_I`.

Remark 3.2 (Prouhet cancellation becomes the normalization). The factor t^{-m} in (3.2) is harmless precisely because P_m has a zero of order m at 1. Thus the same cancellation that annihilates the first m power moments also regularizes the Fourier approximation at the origin.

3.2 Decaying specialization: the endpoint Laplace bridge

Theorem 3.3 (Exact finite Thue–Morse–Laplace identity). *For every integer $m \geq 1$ and every $s \in \mathbb{C}$,*

$$G_m(-s) = \frac{2^{m(m+1)/2}}{s^m} P_m\left(e^{-s/2^m}\right) = \frac{2^{m(m+1)/2}}{s^m} \sum_{n=0}^{2^m-1} \varepsilon_n e^{-ns/2^m}. \quad (3.3)$$

Again, the apparent singularity at $s = 0$ is removable. Result derived in this report.

Proof. The substitution $z = e^{-s/2^m}$ gives

$$P_m(z) = \prod_{j=0}^{m-1} \left(1 - e^{-s2^j/2^m}\right) = \prod_{k=1}^m \left(1 - e^{-s/2^k}\right).$$

The denominator in $G_m(-s)$ is

$$\prod_{k=1}^m \frac{s}{2^k} = s^m 2^{-m(m+1)/2},$$

which proves the identity. $\square$

Formal version. The total denominator-free form underlying (3.3) is `Fabius.thueMorseBlock_exp_neg_eq_laplacePrefix` (`ThueMorseComplexProductBridge.lean`). Its prefix `Fabius.complexLaplacePrefix` is exactly $G_m(-s)$, and the theorem holds for every $m \geq 0$ and every $s \in \mathbb{C}$, including $s = 0$. The normalized origin value is `Fabius.complexLaplacePrefix_apply_zero`. Away from zero, `Fabius.complexLaplacePrefix_eq_thueMorseBlock_exp_neg` is precisely the quotient form printed in (3.3). The factor identity `Fabius.complexExp1Div_neg_eq_exp_mul_complexSinc` and its finite product corollary `Fabius.complexLaplacePrefix_eq_exp_mul_shiftedComplexSincPrefix` also give the exact rotation (5.2), with no condition at the removable origin.

The exact finite-block principle. The oscillatory sample $z = e^{it/2^{m-1}}$ converts P_m into the truncated Fourier product $Q_m(t)$. The decaying sample $z = e^{-s/2^m}$ converts the same P_m into the truncated negative-Laplace product $G_m(-s)$. No asymptotic argument is involved.

3.3 The omitted tail is a scaled copy of the whole

Proposition 3.4 (Exact tail self-similarity). *For every $m \geq 0$,*

$$Q_\infty(t) = Q_m(t)Q_\infty(t/2^m), \quad (3.4)$$

$$G(-s) = G_m(-s)G(-s/2^m). \quad (3.5)$$

Result derived in this report.

Proof. After removing the first m factors, replace k by $m + \ell$ in the remaining infinite product. The residual factors are exactly those of the original product evaluated at the scaled argument. $\square$

Formal versions. Both halves are formalized, and the first is formalized in a form strictly more general than the one stated here. The reindexing itself is `Fabius.tprod_geom_scale_pow` (`GeometricScaleProducts.lean`): for an arbitrary function g from a field into a commutative topological monoid, an arbitrary nonzero scale c , and every k ,

$$\prod_{n \geq 0} g\left(\frac{c^k z}{c^n}\right) = \left(\prod_{j < k} g(c^{j+1} z)\right) \prod_{n \geq 0} g\left(\frac{z}{c^n}\right),$$

the only hypothesis being multipliability of the undilated family; nothing about sinc_{rad} , about base two, or even about $\mathbb{C}$ enters, and $g = \text{sinc}_{\text{rad}}$, $c = 2$, $z = t/2^{m+1}$, $k = m$ gives (3.4). The instance stated directly for the transform is `Fabius.rvachevFourier_eq_finite_product` (`FourierProduct.lean`), in the $e^{-2\pi i x \xi}$ convention of (2.4). The Laplace half (3.5) is `Fabius.complexGeneratingFunction_neg_finite_refinement` (`NegativeLaplaceVertical.lean`), an exact identity for every m and every $s \in \mathbb{C}$.

Corollary 3.5 (Exact finite Thue–Morse factorization of the full products). *For $m \geq 1$,*

$$Q_\infty(t) = i^m 2^{\binom{m}{2}} t^{-m} e^{-it(1-2^{-m})} P_m\left(e^{it/2^{m-1}}\right) Q_\infty(t/2^m), \quad (3.6)$$

$$G(-s) = \frac{2^{m(m+1)/2}}{s^m} P_m\left(e^{-s/2^m}\right) G(-s/2^m). \quad (3.7)$$

Thus every truncation error is represented by the original entire function in a small argument.

4 The zeta–Lambert tail calculus

4.1 A local logarithm of the sinc function

Lemma 4.1 (Euler–zeta expansion). *For $|z| < \pi$,*

$$\log \operatorname{sinc}_{\text{rad}} z = - \sum_{r=1}^{\infty} \frac{\zeta(2r)}{r\pi^{2r}} z^{2r}. \quad (4.1)$$

The logarithm is the branch vanishing at $z = 0$.

Proof. Euler’s product gives

$$\frac{\sin z}{z} = \prod_{n=1}^{\infty} \left(1 - \frac{z^2}{\pi^2 n^2}\right).$$

For $|z| < \pi$, expand $\log(1 - w) = -\sum_{r \geq 1} w^r / r$. Absolute local uniform convergence permits interchanging the sums, and the inner sum over n is $\zeta(2r)$. $\square$

Formal versions. The branch-free content of (4.1) is formalized in `SincZetaSeries.lean` as the exponential identity `Fabius.complexSinc_pi_mul_eq_cexp`: for complex x with $|x| < 1$, $\operatorname{sinc}_{\text{rad}}(\pi x) = \exp(-\sum_{r \geq 1} \zeta(2r)x^{2r}/r)$, the values $\zeta(2r)$ carried as the elementary series `Fabius.evenZeta` (`EvenZetaSeries.lean`); the $|z| < \pi$ normalization printed above is `Fabius.complexSinc_eq_cexp`. For real $0 < |x| < 1$ the boxed logarithmic form itself is `Fabius.log_sin_div_pi_mul`, the branch being the real logarithm – exactly the branch vanishing at 0. The interchange of sums is performed once, in greater generality than this lemma needs (an arbitrary absolutely summable family with every norm below 1, over an arbitrary index type), by the Euler log transform `Fabius.hasSum_powerSum_log_one_sub` of `EulerLogTransform.lean`; the closed evaluations $\zeta(2) = \pi^2/6$ and $\zeta(4) = \pi^4/90$ and the general Bernoulli form are `EvenZeta.Values.lean`.

4.2 All orders, exact radius, explicit remainder

Theorem 4.2 (All-orders sinc-tail expansion). *Let $m \geq 0$. For*

$$|t| < \pi 2^{m+1}, \quad (4.2)$$

the quotient $Q_\infty(t)/Q_m(t)$ is nonzero and

$$\log \frac{Q_\infty(t)}{Q_m(t)} = - \sum_{r=1}^{\infty} \frac{\zeta(2r)}{r\pi^{2r}(4^r - 1)} t^{2r} 4^{-rm}. \quad (4.3)$$

*The disk in (4.2) is maximal for this Taylor logarithm: its boundary contains the nearest zeros $t = \pm \pi 2^{m+1}$ of the tail. **Result derived in this report.***

Proof. By [proposition 3.4](#), the quotient is $Q_\infty(t/2^m)$. Apply [lemma 4.1](#) to every factor and sum:

$$\begin{aligned} \log Q_\infty(t/2^m) &= - \sum_{k=1}^{\infty} \sum_{r=1}^{\infty} \frac{\zeta(2r)}{r\pi^{2r}} \left(\frac{t}{2^{m+k}} \right)^{2r} \\ &= - \sum_{r=1}^{\infty} \frac{\zeta(2r)t^{2r}}{r\pi^{2r}} \frac{4^{-r(m+1)}}{1-4^{-r}} \\ &= - \sum_{r=1}^{\infty} \frac{\zeta(2r)}{r\pi^{2r}(4^r-1)} t^{2r} 4^{-rm}. \end{aligned}$$

The closest zero of $Q_\infty(w)$ to $w = 0$ is $w = \pm 2\pi$, arising from the factor $\text{sinc}_{\text{rad}}(w/2)$. Scaling by $w = t/2^m$ proves maximality. $\square$

Formal versions. The theorem is formalized in `SincZetaDyadic.lean`, in the normalization $\Phi(z) = Q_\infty(2\pi z)$ and in exponential form. At $m = 0$, `Fabius.rvachevFourierProduct_eq_cexp` proves $\Phi(z) = \exp(-\sum_{r \geq 1} \zeta(2r)4^r z^{2r}/(r(4^r-1)))$ on $|z| < 1$; for general m , `Fabius.rvachevFourierProduct_eq_prefix_mul_cexp` is the exact factorization $\Phi(z) = (\prod_{j < m} \text{sinc}_{\text{rad}}(\pi z/2^j)) \exp(-\sum_{r \geq 1} \zeta(2r)4^r (z/2^m)^{2r}/(r(4^r-1)))$ on $|z| < 2^m$ – identity (4.3) with the quotient cleared and the finite prefix explicit. The engine is one application of the Euler log transform to the pair family $(z/2^h)^2/(n+1)^2$ over $\mathbb{N}_0 \times \mathbb{N}_0$; its power sums are evaluated for *every* complex z , with no smallness hypothesis (`Fabius.sinc_pair_powerSum`). Composed with the integral-to-product identity, the same series is the Fourier transform itself: $\widehat{\text{up}}_{2\pi}(z) = \exp(-\sum_{r \geq 1} \zeta(2r)4^r z^{2r}/(r(4^r-1)))$ on $|z| < 1$ (`Fabius.rvachevFourier_eq_cexp`). Maximality of the disk is now formalized too, in `ExponentialDiskMaximality.lean`, and by an obstruction that uses none of the analysis: Φ vanishes at every nonzero integer while $\exp$ never vanishes, so $\Phi(1) = 0$ is not an exponential value at all (`Fabius.rvachevFourierProduct_int_ne_cexp`). Quantifying over the exponent rather than fixing the Euler–zeta series makes the conclusion a statement about the disk rather than about one formula: no disk centred at the origin of radius greater than 1 carries any representation of Φ as an exponential (`Fabius.not_cexpRepresentable_of_one_lt`), and the representable radii have greatest element exactly 1 (`Fabius.isGreatest_cexpRepresentable_radius`). The same argument does *not* settle the prefixed form on $|z| < 2^m$ for $m \geq 1$: there the finite prefix already vanishes at every nonzero integer, its $j = 0$ factor being $\text{sinc}_{\text{rad}}(\pi n)$, so both sides vanish together on the zero set of Φ and the obstruction cannot detect a failure. That too is recorded as a theorem rather than as a sentence (`Fabius.prod_complexSinc_int_eq_zero`), and maximality for $m \geq 1$ stays open.

The first terms are

$$\frac{Q_\infty(t)}{Q_m(t)} = \exp\left(-\frac{t^2}{18}4^{-m} - \frac{t^4}{2700}16^{-m} - \frac{t^6}{178605}64^{-m} - \frac{t^8}{9639000}256^{-m} - \dots\right). \quad (4.4)$$

Theorem 4.3 (Certified tail remainder). *Let $N \geq 0$ and put*

$$x_m(t) = \frac{|t|}{\pi 2^{m+1}} < 1. \quad (4.5)$$

If $R_{N,m}(t)$ denotes the part of the positive series in (4.3) with $r \geq N+1$, then

$$\boxed{|R_{N,m}(t)| \leq \frac{4\zeta(2)}{3} \frac{x_m(t)^{2N+2}}{(N+1)(1-x_m(t)^2)}}. \quad (4.6)$$

For real t in the stated disk, $R_{N,m}(t) \geq 0$. Result derived in this report.

Proof. Rewrite the absolute value of the r -th term as

$$\frac{\zeta(2r)}{r} \frac{4^r}{4^r - 1} x_m(t)^{2r}.$$

Use $\zeta(2r) \leq \zeta(2)$, $4^r/(4^r - 1) \leq 4/3$, and $1/r \leq 1/(N+1)$ for $r \geq N+1$, then sum the geometric series. $\square$

Formal versions. `SincZetaRemainder.lean` packages the coefficient as `Fabius.sincZetaCoeff` and proves (4.6) verbatim as `Fabius.norm_sincZeta_tail_le`, with $x_m(t) = \|w\|$ for $w = z/2^m$, alongside the term-by-term nonnegativity for real arguments (`Fabius.sincZeta_tail_nonneg`) and the three-factor coefficient estimate of this proof (`Fabius.sincZetaCoeff_le`, resting on the antitonicity `Fabius.evenZeta_anti` of the even zeta values). The constant is carried symbolically as $\frac{4}{3}\zeta(2)$; the numerical value $2\pi^2/9$ follows from `Fabius.evenZeta_one` (`EvenZetaValues.lean`).

4.3 Arbitrary-order finite Thue–Morse approximants

Definition 4.4 (Zeta-corrected block approximant). For $N \geq 0$ define

$$Q_m^{[N]}(t) = Q_m(t) \exp \left[- \sum_{r=1}^N \frac{\zeta(2r)t^{2r}}{r\pi^{2r}(4^r - 1)} 4^{-rm} \right]. \quad (4.7)$$

For $N = 0$, this is Q_m .

Corollary 4.5 (Arbitrary-order convergence). For every compact set $K \subset \mathbb{C}$ and every fixed $N \geq 0$,

$$Q_m^{[N]}(t) = Q_\infty(t) \left(1 + \mathcal{O}_K(4^{-(N+1)m}) \right) \quad (t \in K, m \rightarrow \infty). \quad (4.8)$$

Moreover, after inserting (3.2), the left side is a normalized length- 2^m Thue–Morse exponential sum multiplied by only N explicit zeta corrections. **Result derived in this report.**

Proof. For all sufficiently large m , K lies in the disk (4.2). The exact quotient $Q_\infty/Q_m^{[N]}$ is $\exp(-R_{N,m})$. The bound (4.6), uniformly on K , is $\mathcal{O}_K(4^{-(N+1)m})$; exponentiation preserves that order. $\square$

Formal versions. The quantitative content is `Fabius.rvachevFourierProduct_eq_prefix_mul_correction_mul_tail` (`SincZetaRemainder.lean`): on $|z| < 2^m$, exactly $\Phi(z) = \left(\prod_{j < m} \text{sinc}_{\text{rad}}(\pi z/2^j) \right) \cdot \exp\left(-\sum_{r=1}^N c_r(z/2^m)^{2r}\right) \cdot \exp(-R_N(z/2^m))$, the finite prefix being the length- 2^m Thue–Morse block of corollary 3.5, the middle factor the N zeta corrections of (4.7), and the tail exponent bounded by `Fabius.norm_sincZeta_tail_le`. The $\mathcal{O}_K$ formulation printed here is a direct reading of that bound and is not separately formalized.

Corollary 4.6 (Fourier approximation to the up-function). For fixed $\xi \in \mathbb{C}$,

$$\begin{aligned} \widehat{\text{up}}_{2\pi}(\xi) &= i^m 2^{\binom{m}{2}} (2\pi\xi)^{-m} e^{-2\pi i \xi(1-2^{-m})} \sum_{n=0}^{2^m-1} \varepsilon_n e^{2\pi i n \xi/2^{m-1}} \\ &\quad \times \exp \left[- \sum_{r=1}^N \frac{\zeta(2r)(2\pi\xi)^{2r}}{r\pi^{2r}(4^r - 1)} 4^{-rm} \right] \left(1 + \mathcal{O}_{\xi,N}(4^{-(N+1)m}) \right). \end{aligned} \quad (4.9)$$

Thus the infinite Fourier product admits an arbitrary-order approximation from one finite digital block.

4.4 Why the coefficient array is a Lambert series

Put $q = 1/4$. Then (4.3) becomes

$$\log \frac{Q_\infty(t)}{Q_m(t)} = - \sum_{r=1}^{\infty} \frac{\zeta(2r)}{r\pi^{2r}} t^{2r} \frac{q^{r(m+1)}}{1 - q^r}. \quad (4.10)$$

The kernel $q^r/(1 - q^r)$ is the elementary Lambert-series kernel. Expanding it as $\sum_{\ell \geq 1} q^{r\ell}$ reconstructs the omitted factors. This is a direct analytic connection between the sinc product and q -series: the base is $q = 1/4$ because the logarithm is even and the geometric frequency scale is $1/2$.

The kernel itself is now formal, in a generality independent of the sinc product: for complex $|w| < 1$ and $|c| < 1$, `Fabius.tsum_log_one_sub_geom (EulerLogTransform.lean)` proves $\sum_{k \geq 0} \log(1 - wc^k) = - \sum_{r \geq 1} w^r / (r(1 - c^r))$, and `Fabius.tprod_one_sub_geom_eq_cexp` is the exponential q -Pochhammer form $(w; c)_\infty = \exp(- \sum_{r \geq 1} w^r / (r(1 - c^r)))$, with real counterparts (`Fabius.tsum_log_one_sub_geom_real`, `Fabius.tprod_one_sub_geom_eq_rexp`).

Theorem 4.7 (General geometric base). *Let $b > 1$ and*

$$Q_{m,b}(t) = \prod_{k=1}^m \text{sinc}_{\text{rad}}(t/b^k), \quad Q_{\infty,b}(t) = \prod_{k=1}^{\infty} \text{sinc}_{\text{rad}}(t/b^k), \quad q = b^{-2}.$$

For $|t| < \pi b^{m+1}$,

$$\log \frac{Q_{\infty,b}(t)}{Q_{m,b}(t)} = - \sum_{r=1}^{\infty} \frac{\zeta(2r)}{r\pi^{2r}} t^{2r} \frac{q^{r(m+1)}}{1 - q^r}. \quad (4.11)$$

Result derived in this report.

Proof. Repeat the proof of [theorem 4.2](#), replacing the geometric sum in 4^{-r} by the one in b^{-2r} . $\square$

5 An all-orders finite-block formula for the endpoint Laplace product

5.1 Centering removes the linear tail

Lemma 5.1 (Fourier–Laplace rotation). *For every $s \in \mathbb{C}$,*

$$G(-s) = e^{-s/2} Q_\infty(is/2). \quad (5.1)$$

More generally,

$$G_m(-s) = e^{-\frac{s}{2}(1-2^{-m})} Q_m(is/2). \quad (5.2)$$

Proof. For every v ,

$$\frac{1 - e^{-v}}{v} = e^{-v/2} \frac{\sinh(v/2)}{v/2} = e^{-v/2} \text{sinc}_{\text{rad}}(iv/2).$$

Multiply with $v = s/2^k$. The linear exponents sum to $-s/2$ in the infinite product and to $-s(1 - 2^{-m})/2$ in the finite product. $\square$

Formal version. The infinite case (5.1) is formalized, as the composite of three named theorems. `Fabius.complexGeneratingFunction_neg_eq_tprod (NegativeLaplaceVertical.lean)` identifies $G(-s)$ with the entire half-moment generating function, as recorded after (2.3); `Fabius.complexGeneratingFunction_eq_fourier_analytic (FourierAnalytic.lean)` states that generating function to be $e^{z/2} \widehat{\text{up}}_{2\pi}(iz/(4\pi))$ for every $z \in \mathbb{C}$; and `Fabius.rvachev_Fourier_eq_product (FourierProduct.lean)` replaces $\mathcal{F}_{2\pi}[\text{up}]$ by its sinc product. Since

$\widehat{\text{up}}_{2\pi}(iz/(4\pi)) = Q_\infty(iz/2)$ in the normalization (2.4), and Q_∞ is even, the composite at $z = -s$ is exactly (5.1). The finite form (5.2) is the total theorem `Fabius.complexLaplacePrefix_eq_exp_mul_shiftedComplexSincPrefix` in `ThueMorseComplexProductBridge.lean`; its factor-level engine is `Fabius.complexExpn1Div_neg_eq_exp_mul_complexSinc`.

Define the centered products

$$C(-s) = e^{s/2} G(-s), \quad C_m(-s) = e^{\frac{s}{2}(1-2^{-m})} G_m(-s). \quad (5.3)$$

Then $C(-s) = Q_\infty(is/2)$ and

$$C(-s) = C_m(-s) C(-s/2^m). \quad (5.4)$$

The centering is not cosmetic: the raw tail has a first-order error, while the centered tail has a second-order error.

5.2 Bernoulli–zeta expansion of the negative tail

Theorem 5.2 (All-orders negative-Laplace tail). *For $|u| < 4\pi$,*

$$\log G(-u) = -\frac{u}{2} + \sum_{r=1}^{\infty} \frac{B_{2r}}{2^r (2r)! (4^r - 1)} u^{2r}. \quad (5.5)$$

Equivalently,

$$\log C(-u) = \frac{u^2}{72} - \frac{u^4}{43200} + \frac{u^6}{11430720} - \frac{u^8}{2467584000} + \cdots \quad (5.6)$$

The radius 4π is exact. **Result derived in this report.**

Proof. Use lemma 5.1 and theorem 4.2 with $m = 0$ and $t = iu/2$:

$$\log G(-u) = -\frac{u}{2} - \sum_{r \geq 1} \frac{\zeta(2r)}{r \pi^{2r} (4^r - 1)} (-1)^r \frac{u^{2r}}{4^r}.$$

The standard evaluation

$$\zeta(2r) = (-1)^{r+1} \frac{B_{2r} (2\pi)^{2r}}{2(2r)!}$$

reduces the coefficient to the one displayed. The closest zeros of $G(-u)$ are $u = \pm 4\pi i$, coming from $1 - e^{-u/2}$, so the logarithmic radius is exact. $\square$

Corollary 5.3 (Certified remainder for the centered Laplace tail). *Let $y = |u|/(4\pi) < 1$. If the series in (5.5) is stopped after u^{2N} , while the exact linear term $-u/2$ is retained, then the remainder $\mathcal{R}_N(u)$ obeys*

$$|\mathcal{R}_N(u)| \leq \frac{4\zeta(2)}{3} \frac{y^{2N+2}}{(N+1)(1-y^2)}. \quad (5.7)$$

Proof. This is theorem 4.3 at $m = 0$, $t = iu/2$. $\square$

5.3 The endpoint product from one finite digital block

Theorem 5.4 (Finite Thue–Morse endpoint formula with all corrections). *Let $m \geq 1$, $s \in \mathbb{C}$, and $u = s/2^m$. If $|u| < 4\pi$, then*

$$\begin{aligned} G(-s) &= \frac{2^{m(m+1)/2}}{s^m} P_m(e^{-u}) \\ &\quad \times \exp \left[-\frac{u}{2} + \sum_{r=1}^{\infty} \frac{B_{2r}}{2^r (2r)! (4^r - 1)} u^{2r} \right]. \end{aligned} \quad (5.8)$$

Stopping the correction series after $r = N$ incurs the explicit logarithmic error (5.7). In sum form, $P_m(e^{-u}) = \sum_{n < 2^m} \varepsilon_n e^{-nu}$. **Result derived in this report.**

Proof. Combine the exact factorization (3.7) with theorem 5.2 applied to the tail $G(-u)$. $\square$

Remark 5.5 (Arithmetic-to-asymptotic bridge). The exact dyadic arithmetic literature typically uses finite coefficient extraction and q -binomial identities, whereas endpoint asymptotics use the infinite product $G(-s)$ and a saddle. Formula (5.8) gives an intermediate object: a single finite Thue–Morse exponential sum, multiplied by a convergent universal correction whose coefficients are Bernoulli or zeta values. It is exact within its natural logarithmic disk, not merely asymptotic in m .

6 q -binomial Richardson acceleration on the spectral scale

6.1 An analytic error germ in q^m

Put $q = 1/4$. From (3.4),

$$Q_m(t) = Q_\infty(t)H_t(q^m), \quad H_t(r) = \frac{1}{Q_\infty(t\sqrt{r})}. \quad (6.1)$$

Although a square root appears in the notation, Q_∞ is even, so H_t is a genuine meromorphic function of r . Its Taylor expansion at 0 is

$$H_t(r) = \exp\left(\sum_{\nu \geq 1} \frac{\zeta(2\nu)t^{2\nu}}{\nu\pi^{2\nu}(4^\nu - 1)} r^\nu\right) = \sum_{n \geq 0} h_n(t)r^n. \quad (6.2)$$

For $t \neq 0$, the nearest pole occurs at $r = 4\pi^2/t^2$. The first coefficients are

$$h_0 = 1, \quad h_1 = \frac{t^2}{18}, \quad h_2 = \frac{31t^4}{16200}, \quad h_3 = \frac{2347t^6}{42865200}. \quad (6.3)$$

6.2 Geometric-grid weights are Gaussian binomials

For $(q; q)_n = \prod_{k=1}^n (1 - q^k)$, define

$$w_{p,j} = \frac{(-1)^{p-j} q^{\binom{p-j+1}{2}}}{(q; q)_j (q; q)_{p-j}} = \frac{(-1)^{p-j} q^{\binom{p-j+1}{2}}}{(q; q)_p} \begin{bmatrix} p \\ j \end{bmatrix}_q, \quad 0 \leq j \leq p. \quad (6.4)$$

Lemma 6.1 (Lagrange characterization). *The numbers in (6.4) are the Lagrange weights for evaluating at $r = 0$ from the geometric nodes $1, q, \dots, q^p$:*

$$w_{p,j} = \prod_{\substack{0 \leq \ell \leq p \\ \ell \neq j}} \frac{-q^\ell}{q^j - q^\ell}. \quad (6.5)$$

Consequently,

$$\sum_{j=0}^p w_{p,j} = 1, \quad \sum_{j=0}^p w_{p,j} q^{jn} = 0 \quad (1 \leq n \leq p), \quad (6.6)$$

and

$$\sum_{j=0}^p w_{p,j} q^{j(p+1)} = (-1)^p q^{p(p+1)/2}. \quad (6.7)$$

Proof. Split the denominator in (6.5) into $\ell < j$ and $\ell > j$, factor the powers of q , and collect the factors $1 - q^k$. This gives (6.4). The moment identities state that Lagrange interpolation reproduces the monomials $1, r, \dots, r^p$ at $r = 0$. For r^{p+1} , the interpolation error is the nodal polynomial $\prod_{\ell=0}^p (r - q^\ell)$; evaluation at 0 gives $(-1)^{p+1} q^{p(p+1)/2}$ for $f - P$, hence (6.7) for $P - f$. $\square$

Formal versions. The interpolation characterization and all three moment identities are formalized in `GeometricLagrange.lean`, over an arbitrary field and for an arbitrary ratio q – not merely at $q = 1/4$. The Lean weight `Fabius.geometricLagrangeWeight` is *defined* as the evaluation-at-zero Lagrange weight for the nodes $1, q, \dots, q^n$, that is, by (6.5); the distinctness of the nodes is not built into the definition but carried as an explicit injectivity hypothesis on $j \mapsto q^j$. With that hypothesis, `Fabius.sum_geometricLagrangeWeight` is the normalization $\sum_j w_{p,j} = 1$, `Fabius.sum_geometricLagrangeWeight_mul_pow_eq_zero` is the vanishing $\sum_j w_{p,j} q^{jn} = 0$ for $1 \leq n \leq p$ – together (6.6) – and `Fabius.sum_geometricLagrangeWeight_mul_pow_succ` is (6.7), with the product form of the same defect recorded separately as `Fabius.sum_geometricLagrangeWeight_mul_pow_succ_product`. The closed form (6.4) itself – the passage from the Lagrange quotient to the q -Pochhammer and Gaussian-binomial expression – is formalized in the companion module described next.

Formal crosswalk. The compiled module `GeometricQBinomialLagrange.lean` packages this finite algebra at arbitrary ratio. Write $N_{p,j}(q)$ for `geometricQBinomialWeightNumerator` applied to q, p, j . It is zero for $j > p$, and for $j \leq p$ is definitionally given in the reversed form $(-1)^{p-j} q^{\binom{p-j+1}{2}} \text{gaussianBinomial } q \, p \, (p-j)$. The division-free symmetry theorem `gaussianBinomial_symm` in `FiniteQBinomialCore.lean` and zero extension above the diagonal now give the all-index forward formula

$$N_{p,j}(q) = (-1)^{p-j} q^{\binom{p-j+1}{2}} \text{gaussianBinomial } q \, p \, j$$

as `geometricQBinomialWeightNumerator_eq_forward`; thus the displayed coefficient $\begin{bmatrix} p \\ j \end{bmatrix}_q$ in (6.4) is now the literal proved forward-index form. The same core also proves `map_gaussianBinomial`, both q -Pascal recurrences, the q -Pochhammer concatenation laws, the full denominator-free product–Gaussian factorial identity, and the finite q -binomial theorem. The theorem `reversed_finite_qBinomial_theorem` proves the denominator-free polynomial identity with product $\prod_{k < p} (z - q^{k+1})$. Under injectivity of $k \mapsto q^k$ on $\{0, \dots, p\}$, `geometricLagrangeWeight_eq_gaussianBinomial_div` identifies the normalized numerator with $w_{p,j}$, and `sum_geometricLagrangeWeight_mul_pow_eq_gaussianBinomial` gives every positive moment:

$$\sum_{j=0}^p w_{p,j} (q^j)^d = (-1)^p q^{p(p+1)/2} \text{gaussianBinomial } q \, (d-1) \, p \quad (d > 0).$$

This discharges the algebraic weight and residual-moment obligations, including $q = 1/4$ whenever the nodes are injective. At $q = 1/2$, the existing `discreteLimitWeight_eq_geometricLagrangeWeight` is the half-base Toeplitz specialization.

The coefficientwise action of the filter has also been factored at its natural level. Over any commutative semiring, `FinitePowerSeriesFilter.lean` proves

$$[X^m] \sum_{j \in S} w_j f(a_j X) = \left(\sum_{j \in S} w_j a_j^m \right) [X^m] f,$$

together with constant-term, exactness, rescaling, and coefficient-map laws. `GeometricPowerSeriesFilter.lean` specializes this to $a_j = q^j$: `geometricLagrangeSeriesFilter_`

eq_residual_mk packages degree zero and every positive Gaussian residual multiplier in one formal-series identity, while map_geometricSeriesFilter transports the entire identity along a commutative-semiring homomorphism. These are identities in PowerSeries. The new AnalyticSeriesFilter.lean separately proves the direct convergent-series analogue: under unconditional summability of each contributing sampled series it interchanges the finite node sum with the infinite series and gives the exact Gaussian residual tail, including the shifted geometric-Lagrange row. Its hypotheses are imposed only at nonzero weights, so a zero-weight node requires no convergence of the unweighted sample. It neither covers conditionally-only boundary series nor identifies the abstract coefficients with the functions Q_m ; it also supplies no sign, bound, asymptotic, or uniform-convergence theorem. Separately, RvachevQBinomialFilter.lean uses the named even modes from AnalyticMoments.lean to instantiate the filter for the actual infinite Rvachev product, with only finite-power injectivity in the generic complex theorem and no hypothesis at $c = 1/2$. It does not identify the finite prefixes Q_m in theorem 6.2; that finite-prefix bridge and the uniform analytic remainder claims remain frontier.

Theorem 6.2 (q -binomial spectral Richardson extrapolation). *Define*

$$\mathcal{R}_{m,p}(t) = \sum_{j=0}^p w_{p,j} Q_{m+j}(t), \quad q = \frac{1}{4}. \quad (6.8)$$

For every compact $K \subset \mathbb{C}$ and fixed p ,

$$\boxed{\mathcal{R}_{m,p}(t) = Q_\infty(t) + \mathcal{O}_K(q^{(p+1)m}) \quad (t \in K).} \quad (6.9)$$

More precisely,

$$\mathcal{R}_{m,p}(t) - Q_\infty(t) = Q_\infty(t)(-1)^p q^{p(p+1)/2} h_{p+1}(t) q^{(p+1)m} + \mathcal{O}_K(q^{(p+2)m}). \quad (6.10)$$

Every Q_{m+j} may be replaced by its exact finite Thue–Morse representation (3.2); hence (6.8) is a finite digital extrapolation formula with Gaussian-binomial coefficients. **Result derived in this report.**

Proof. Insert (6.2) into (6.8):

$$\begin{aligned} \frac{\mathcal{R}_{m,p}(t)}{Q_\infty(t)} &= \sum_{j=0}^p w_{p,j} \sum_{n \geq 0} h_n(t) q^{n(m+j)} \\ &= \sum_{n \geq 0} h_n(t) q^{nm} \left(\sum_{j=0}^p w_{p,j} q^{jn} \right). \end{aligned}$$

The inner sums are 1 at $n = 0$, vanish for $1 \leq n \leq p$, and equal (6.7) at $n = p + 1$. Local uniform convergence of the analytic H_t series gives the uniform remainder on compact t -sets. $\square$

Corollary 6.3 (A derivative bound version). *Suppose $t \in \mathbb{R}$ and the interval $[0, q^m]$ contains no pole of H_t . Then for some $\theta \in (0, q^m)$,*

$$\frac{\mathcal{R}_{m,p}(t)}{Q_\infty(t)} - 1 = \frac{(-1)^p q^{(p+1)m+p(p+1)/2}}{(p+1)!} H_t^{(p+1)}(\theta). \quad (6.11)$$

In particular, a supremum bound on $H_t^{(p+1)}$ gives a fully certified extrapolation error.

Proof. Apply the real Lagrange interpolation remainder to H_t at the nodes $q^m, q^{m+1}, \dots, q^{m+p}$ and evaluate the interpolating polynomial at zero. $\square$

The first four weight rows are

$$\begin{aligned}
 p = 1 : & \quad \left(-\frac{1}{3}, \frac{4}{3} \right), \\
 p = 2 : & \quad \left(\frac{1}{45}, -\frac{4}{9}, \frac{64}{45} \right), \\
 p = 3 : & \quad \left(-\frac{1}{2835}, \frac{4}{135}, -\frac{64}{135}, \frac{4096}{2835} \right), \\
 p = 4 : & \quad \left(\frac{1}{722925}, -\frac{4}{8505}, \frac{64}{2025}, -\frac{4096}{8505}, \frac{1048576}{722925} \right).
 \end{aligned} \tag{6.12}$$

Remark 6.4 (Two complementary accelerations). The zeta correction (4.7) uses one block level and analytic coefficients. The q -Richardson formula (6.8) uses $p + 1$ consecutive block levels but no zeta values. Both exploit the same analytic germ in 4^{-m} , so their observed orders agree: correction through N and extrapolation degree $p = N$ both leave order $4^{-(N+1)m}$.

7 Exact random tails and a weak differential expansion

7.1 The random-series model

Let $V_1, V_2, \dots$ be independent random variables uniformly distributed on $[-1, 1]$, and put

$$Y_m = \sum_{k=1}^m 2^{-k} V_k, \quad Y_\infty = \sum_{k=1}^{\infty} 2^{-k} V_k. \tag{7.1}$$

Then Y_∞ has density up , while Y_m has the compactly supported centered uniform spline density u_m . Their characteristic functions are Q_∞ and Q_m .

Theorem 7.1 (Exact random-tail decomposition). *There is an independent copy Y'_∞ of Y_∞ such that*

$$\boxed{Y_\infty \stackrel{d}{=} Y_m + 2^{-m} Y'_\infty.} \tag{7.2}$$

Consequently, in density form,

$$\boxed{\text{up} = u_m * (2^m \text{up}(2^m \cdot)).} \tag{7.3}$$

Result derived in this report.

Proof. Split the infinite sum after m and reindex:

$$\sum_{k=m+1}^{\infty} 2^{-k} V_k = 2^{-m} \sum_{\ell=1}^{\infty} 2^{-\ell} V_{m+\ell}.$$

The reindexed series is independent of Y_m and has the law of Y_∞ . Convolution of the corresponding densities gives (7.3). $\square$

Formal versions. The measure form of (7.3) is formalized in `MeasureRefinement.lean`: with μ_{up} the Borel probability measure with density up (`Fabius.rvachevMeasure`), the one-step refinement equation $\mu_{\text{up}} = \text{Unif}[-\frac{1}{2}, \frac{1}{2}] * (\mu_{\text{up}} \circ (\cdot/2)^{-1})$ is `Fabius.rvachevMeasure_refinement`, and the m -fold form matching the split in this proof — the first m uniform digits convolved with a 2^{-m} -scale copy of the whole — is `Fabius.rvachevMeasure_eq_prefix_conv`, with the digit-prefix law `Fabius.uniformDigitPrefix` defined by the digit recursion and its characteristic function evaluated as the finite sinc product (`Fabius.charFun_uniformDigitPrefix`); the uniform-digit integral itself is the reusable `Fabius.charFun_volume_restrict_Icc`. The companion

`UniformDigitThueMorseBridge.lean` identifies that characteristic function with the normalized prefix in `Fabius.charFun_uniformDigitPrefix_eq_shiftedComplexSincPrefix`, substitutes it into the total finite block theorem in `Fabius.thueMorseBlock_cexp_eq_charFun_uniformDigitPrefix`, and solves back for the characteristic function away from zero in `Fabius.charFun_uniformDigitPrefix_eq_thueMorseBlock_cexp`. These are finite-prefix identities; they assert no infinite-product convergence or random-tail limit. The measure-refinement proof runs by exactly the characteristic-function route anticipated in the promotion table below. On the random-variable side, the representation itself is formalized independently in `ProbabilityRepresentation.lean`: the product probability space of uniform digits, the uncentered random series $X_\infty = (Y_\infty + 1)/2$, and the identifications $F(x) = \mathbb{P}[X_\infty \leq x]$ and $\text{up}(x) = \mathbb{P}[X_\infty \leq 1 - |x|]$ (Arias de Reyna’s Theorem 3). The bridge between the two sides is formalized in `RandomSeriesLaw.lean`: the up-measure is the pushforward of the law of X_∞ under $x \mapsto 2x - 1$ (`Fabius.rvachevMeasure_eq_map_weightedSum`, proved by the fundamental theorem of calculus for the folded derivative $F' = 2 \text{up}(2 \cdot - 1)$, which exhibits $y \mapsto F(\frac{y+1}{2})$ as the CDF of μ_{up}); through it the characteristic function of the uncentered series takes the closed form $\mathbb{E}[e^{itX_\infty}] = e^{it/2} \widehat{\text{up}}_{2\pi}(t/(4\pi))$ (`Fabius.charFun_weightedSumDistribution`). The one-digit equality-in-law splitting underlying (7.2) is the law identity $\text{law}(X_\infty) = (\text{Unif}[0, 1] * \text{law}(X_\infty)) \circ (\cdot/2)^{-1}$ (`Fabius.weightedSumDistribution_self_similar`). Together with the m -fold prefix-convolution theorem above, this supplies the measure/equality-in-law content of (7.2) for every m ; no separate existential coupling theorem is claimed.

7.2 Cumulants are the same zeta–Lambert coefficients

Let κ_j and μ_j denote the cumulants and moments of Y_∞ .

Theorem 7.2 (Closed cumulant formula). *All odd cumulants vanish, and for $r \geq 1$,*

$$\kappa_{2r} = (-1)^{r+1} \frac{(2r)! \zeta(2r)}{r \pi^{2r} (4^r - 1)} = \frac{2^{2r-1} B_{2r}}{r(4^r - 1)}. \quad (7.4)$$

In particular,

$$\kappa_2 = \frac{1}{9}, \quad \kappa_4 = -\frac{2}{225}, \quad \kappa_6 = \frac{16}{3969}, \quad \kappa_8 = -\frac{16}{3825}. \quad (7.5)$$

Result derived in this report.

Proof. The cumulant expansion of the characteristic function is

$$\log \mathbb{E} e^{itY_\infty} = \sum_{r \geq 1} \kappa_{2r} \frac{(it)^{2r}}{(2r)!}.$$

Set $m = 0$ in (4.3) and identify coefficients. The Bernoulli form follows from the classical evaluation of $\zeta(2r)$. $\square$

The first even moments, obtained from the complete Bell polynomials in the cumulants, are

$$\mu_2 = \frac{1}{9}, \quad \mu_4 = \frac{19}{675}, \quad \mu_6 = \frac{583}{59535}, \quad \mu_8 = \frac{132809}{32531625}. \quad (7.6)$$

7.3 A certified weak expansion

Theorem 7.3 (Weak differential expansion of the up-law). *Let $N \geq 0$ and $\varphi \in C^{2N+2}([-1, 1])$. Then*

$$\mathbb{E} \varphi(Y_\infty) = \sum_{j=0}^N \frac{\mu_{2j}}{(2j)!} 4^{-jm} \mathbb{E} \varphi^{(2j)}(Y_m) + \mathcal{E}_{N,m}(\varphi), \quad (7.7)$$

where $\mu_0 = 1$ and

$$|\mathcal{E}_{N,m}(\varphi)| \leq \frac{4^{-(N+1)m}}{(2N+2)!} \|\varphi^{(2N+2)}\|_{L^\infty([-1,1])}. \quad (7.8)$$

Result derived in this report.

Proof. Condition on Y_m in (7.2) and Taylor-expand $\varphi(Y_m + 2^{-m}Y'_\infty)$ through degree $2N + 1$. Symmetry kills every odd moment. Since $|Y'_\infty| \leq 1$, the Taylor remainder is bounded by $2^{-(2N+2)m} \|\varphi^{(2N+2)}\|_\infty / (2N+2)!$. $\square$

The first terms are therefore

$$\begin{aligned} \mathbb{E}\varphi(Y_\infty) &= \mathbb{E}\varphi(Y_m) + \frac{4^{-m}}{18} \mathbb{E}\varphi''(Y_m) + \frac{19 \cdot 16^{-m}}{16200} \mathbb{E}\varphi^{(4)}(Y_m) \\ &\quad + \frac{583 \cdot 64^{-m}}{42865200} \mathbb{E}\varphi^{(6)}(Y_m) + \mathcal{O}_\varphi(256^{-m}). \end{aligned} \quad (7.9)$$

Corollary 7.4 (Distributional spline expansion). *In the sense of distributions, for every fixed N ,*

$$\text{up} = \sum_{j=0}^N \frac{\mu_{2j}}{(2j)!} 4^{-jm} D^{2j} u_m + \mathcal{O}_{\mathcal{D}'}(4^{-(N+1)m}). \quad (7.10)$$

In particular,

$$\text{up} = u_m + \frac{4^{-m}}{18} D^2 u_m + \frac{19 \cdot 16^{-m}}{16200} D^4 u_m + \dots \quad (7.11)$$

Proof. Pair the asserted distribution with a test function and use $\langle D^{2j} u_m, \varphi \rangle = \mathbb{E}\varphi^{(2j)}(Y_m)$. $\square$

Remark 7.5 (Cumulant operator). Formally, and exactly at the level of the local Fourier logarithm,

$$\text{up} = \exp\left(\sum_{r \geq 1} \frac{\kappa_{2r}}{(2r)!} 4^{-rm} D^{2r}\right) u_m. \quad (7.12)$$

Expanding the exponential changes cumulants into moments and recovers (7.10). This is a deconvolution calculus for the finite uniform splines already present in the repository's repeated-integration and probabilistic constructions.

8 Lambert- W complexity of endpoint product truncation

8.1 The canonical scale and the tail variable

For $0 < x < 1/(eL)$, let $\lambda(x)$ and $s_0(x)$ be given by (2.6)–(2.7). At an integer truncation level m , the entire omitted product is governed by the single tail variable

$$u_{x,m} = \frac{s_0(x)}{2^m} = 2^{\lambda(x)-m}. \quad (8.1)$$

If $m = \lceil \lambda(x) + h \rceil$, then

$$2^{-h-1} < u_{x,m} \leq 2^{-h}. \quad (8.2)$$

Thus each extra factor beyond the Lambert phase halves the raw tail variable.

Theorem 8.1 (Raw and centered depth laws). *Let*

$$\delta_{\text{raw}}(u) = |G(-u) - 1|, \quad (8.3)$$

$$\delta_{\text{ctr}}(u) = |e^{u/2}G(-u) - 1| = |C(-u) - 1|. \quad (8.4)$$

As $u \rightarrow 0^+$,

$$\delta_{\text{raw}}(u) = \frac{u}{2} + \mathcal{O}(u^2), \quad \delta_{\text{ctr}}(u) = \frac{u^2}{72} + \mathcal{O}(u^4). \quad (8.5)$$

Consequently, for $m = \lceil \lambda(x) + h \rceil$ and $h \rightarrow \infty$,

$$\delta_{\text{raw}}(u_{x,m}) = \Theta(2^{-h}), \quad \delta_{\text{ctr}}(u_{x,m}) = \Theta(4^{-h}), \quad (8.6)$$

with absolute two-sided constants once h is sufficiently large. If $m_{\text{raw}}(x, \eta)$ and $m_{\text{ctr}}(x, \eta)$ denote the least depths which make the respective errors at most η , then as $\eta \downarrow 0$,

$$m_{\text{raw}}(x, \eta) = -\frac{W_{-1}(-Lx)}{L} + \log_2 \frac{1}{\eta} + \mathcal{O}(1), \quad (8.7)$$

$$m_{\text{ctr}}(x, \eta) = -\frac{W_{-1}(-Lx)}{L} + \frac{1}{2} \log_2 \frac{1}{\eta} + \mathcal{O}(1). \quad (8.8)$$

The $\mathcal{O}(1)$ constants may be chosen uniformly for x in any sufficiently small interval $(0, x_0]$. **Result derived in this report.**

Proof. Equation (8.5) is the beginning of (5.5):

$$G(-u) = \exp\left(-\frac{u}{2} + \frac{u^2}{72} + \mathcal{O}(u^4)\right), \quad C(-u) = \exp\left(\frac{u^2}{72} + \mathcal{O}(u^4)\right).$$

Hence there are positive constants c_i, C_i and u_0 such that $c_1 u \leq \delta_{\text{raw}}(u) \leq C_1 u$ and $c_2 u^2 \leq \delta_{\text{ctr}}(u) \leq C_2 u^2$ for $0 < u < u_0$. Combine these inequalities with (8.2). Solving $2^{-h} \asymp \eta$ and $4^{-h} \asymp \eta$, then adding the base depth $\lambda(x)$ and absorbing the ceiling into $\mathcal{O}(1)$, gives (8.7)–(8.8). $\square$

Remark 8.2 (Why centering halves the precision surcharge). At the leading saddle scale, reaching relative tolerance η in the raw product requires roughly $\log_2(1/\eta)$ factors beyond λ . Removing the exactly known linear factor $e^{-u/2}$ makes the residual even, so only half as many additional factors are required. This is a concrete algorithmic consequence of the probability centering hidden in lemma 5.1.

8.2 A finite-block endpoint evaluation protocol

Algorithm 8.3 (Lambert-aware finite-block product evaluation). Given small $x > 0$, target logarithmic correction order N , and tolerance η :

1. Compute $\lambda = -W_{-1}(-x \log 2) / \log 2$ and $s_0 = 2^\lambda = \lambda/x$.
2. Choose

$$m = \left\lceil \lambda + \frac{1}{2} \log_2 \frac{1}{\eta} + c_N \right\rceil,$$

where c_N is a small safety margin determined from (5.7).

3. Set $u = s_0/2^m$ and evaluate the finite digital sum

$$P_m(e^{-u}) = \sum_{n=0}^{2^m-1} \varepsilon_n e^{-nu}$$

or its product form.

4. Reconstruct $G(-s_0)$ from (5.8), using the first N Bernoulli corrections and the explicit remainder bound.

Warning 8.4. This protocol certifies the product evaluation at the leading Lambert scale. Computing $F(x)$ itself still requires the repository’s saddle inversion, including the nonconstant periodic phase and its derivatives. The result here supplies a new certified inner engine for that larger endpoint calculation; it does not suppress the periodic phase.

9 A Thue–Morse Barnes-type tower from Dirichlet trivial zeros

9.1 Input from the Atlas

The Thue–Morse Atlas establishes that, for $a > 0$, the shifted Dirichlet series

$$D(s, a) = \sum_{n \geq 0} \frac{\varepsilon_n}{(n + a)^s} \quad (9.1)$$

has an entire continuation $\tilde{D}(s, a)$, satisfies

$$\tilde{D}(s, a) = 2^{-s} \left[\tilde{D}\left(s, \frac{a}{2}\right) - \tilde{D}\left(s, \frac{a+1}{2}\right) \right], \quad (9.2)$$

and has trivial zeros

$$\tilde{D}(-r, a) = 0 \quad (r = 0, 1, 2, \dots). \quad (9.3)$$

It also supplies the Mellin representation

$$\tilde{D}(s, a) = \frac{1}{\Gamma(s)} \int_0^\infty t^{s-1} e^{-at} \mathcal{E}(t) dt, \quad \mathcal{E}(t) = \prod_{j \geq 0} (1 - e^{-2^j t}), \quad (9.4)$$

where the extraordinary flatness of $\mathcal{E}$ at zero makes the continuation entire [3, 9].

Formal versions of the input. All three Atlas inputs used in this section are formalized, in `ThueMorseEntireContinuation.lean`. The Mellin representation (9.4) is not a theorem there but the *definition* of the continuation: `Fabius.dirichletMellinContinuation` is $s \mapsto \Gamma(s)^{-1} \int_0^\infty t^{s-1} e^{-at} \mathcal{E}(t) dt$, with $\mathcal{E}$ the lacunary product `Fabius.lacunaryExpProduct` and the kernel $e^{-at} \mathcal{E}(t)$ the object `Fabius.mellinKernel`. Entirety is `Fabius.dirichletMellinContinuation_differentiable`, resting on `Fabius.mellin_mellinKernel_differentiable`, which is where the boundary flatness of $\mathcal{E}$ is actually spent. The dyadic equation (9.2) is `Fabius.dirichletMellinContinuation_dyadic`, proved for every $a > 0$ and every $s \in \mathbb{C}$ by analytic continuation from the ray $\sigma > 1$, so it is available at the nonpositive integers where this section differentiates it. The trivial zeros (9.3) are `Fabius.dirichletMellinContinuation_neg_natCast`; as that declaration’s own documentation stresses, in isolation it is only the formal statement $\Gamma(-r)^{-1} = 0$, and the mathematical assertion – that the *entire* function vanishes at $-r$ – is carried jointly with the differentiability of the Mellin factor, so the two should be cited as a pair.

9.2 The derivative tower

Definition 9.1 (Thue–Morse gamma tower). For each integer $r \geq 0$ and $a > 0$, define

$$\Gamma_r^{\text{TM}}(a) = \exp\left(\partial_s \tilde{D}(-r, a)\right). \quad (9.5)$$

Theorem 9.2 (Mellin, dyadic, and differential laws). *The functions (9.5) satisfy:*

$$\log \Gamma_r^{\text{TM}}(a) = (-1)^r r! \int_0^\infty t^{-r-1} e^{-at} \mathcal{E}(t) dt, \quad (9.6)$$

$$\Gamma_r^{\text{TM}}(a) = \left(\frac{\Gamma_r^{\text{TM}}(a/2)}{\Gamma_r^{\text{TM}}((a+1)/2)} \right)^{2^r}, \quad (9.7)$$

$$\frac{d}{da} \log \Gamma_r^{\text{TM}}(a) = r \log \Gamma_{r-1}^{\text{TM}}(a) \quad (r \geq 1). \quad (9.8)$$

More generally, for $0 \leq k \leq r$,

$$\frac{d^k}{da^k} \log \Gamma_r^{\text{TM}}(a) = (r)^k \log \Gamma_{r-k}^{\text{TM}}(a). \quad (9.9)$$

Result derived in this report.

Proof. Near $s = -r$, the reciprocal gamma function has the expansion

$$\frac{1}{\Gamma(s)} = (-1)^r r! (s + r) + \mathcal{O}((s + r)^2).$$

The Mellin integral in (9.4) is finite there, so differentiation gives (9.6).

Differentiate (9.2) with respect to s and set $s = -r$. The derivative of 2^{-s} is multiplied by the bracket of values at $s = -r$, which vanishes by (9.3). Hence

$$\partial_s \tilde{D}(-r, a) = 2^r \left[\partial_s \tilde{D}\left(-r, \frac{a}{2}\right) - \partial_s \tilde{D}\left(-r, \frac{a+1}{2}\right) \right].$$

Exponentiation gives (9.7).

Finally, differentiation of the Dirichlet/Mellin continuation in the parameter gives

$$\partial_a \tilde{D}(s, a) = -s \tilde{D}(s + 1, a).$$

Differentiate this identity in s :

$$\partial_a \partial_s \tilde{D}(s, a) = -\tilde{D}(s + 1, a) - s \partial_s \tilde{D}(s + 1, a).$$

At $s = -r$ with $r \geq 1$, the first term vanishes by (9.3), leaving $r \partial_s \tilde{D}(1 - r, a)$. This proves (9.8); iteration gives (9.9). $\square$

Corollary 9.3 (The Woods–Robbins layer). *If $G_{\text{WR}}(a, b)$ denotes the convergent Thue–Morse product used in the Atlas, for which*

$$\log G_{\text{WR}}(a, b) = \partial_s \tilde{D}(0, b) - \partial_s \tilde{D}(0, a),$$

then

$$G_{\text{WR}}(a, b) = \frac{\Gamma_0^{\text{TM}}(b)}{\Gamma_0^{\text{TM}}(a)}. \quad (9.10)$$

Thus the known product is the ratio law at the base of an infinite dyadic differential tower.

Formal version and exact current API. The older origin-only bridge remains `Fabius.mpLimit_eq_deriv_sub` in `ThueMorseGDirichlet.lean`, for all $a, b > 0$, with `Fabius.mpLimit` the limit of $\sum_n \varepsilon_n \log((n+a)/(n+b))$. The post-checkpoint source commit 0ba35abd4 closes every spectral level used in the Mellin and dyadic parts of [theorem 9.2](#). First, `ReciprocalGammaJets.lean` has exactly five public declarations: `Fabius.deriv_Gamma_inv_neg_nat`, `Fabius.hasDerivAt_Gamma_inv_neg_nat`, `Fabius.hasDerivAt_Gamma_inv_zero`, `Fabius.analyticOrderAt_Gamma_inv_neg_nat`, and `Fabius.tendsto_Gamma_inv_div_add_nat`. For every $r \in \mathbb{N}_0$, these give

$$(\Gamma^{-1})'(-r) = (-1)^r r!, \quad \text{ord}_{-r}(\Gamma^{-1}) = 1, \quad \frac{\Gamma(s)^{-1}}{s+r} \longrightarrow (-1)^r r! \quad (s \rightarrow -r, s \neq -r).$$

They concern the entire reciprocal function $s \mapsto \Gamma(s)^{-1}$; no derivative value is assigned to the raw Gamma function at a pole.

At the same commit, `ThueMorseGammaTower.lean` has exactly eight public declarations: `Fabius.hasDerivAt_dirichletMellinContinuation_neg_nat`, `Fabius.deriv_dirichletMellinContinuation_neg_nat`, `Fabius.thueMorseGammaLog`, `Fabius.thueMorseGammaTower`, `Fabius.thueMorseGammaLog_eq_mellin`, `Fabius.thueMorseGammaLog_eq_integral`, `Fabius.thueMorseGammaLog_dyadic`, and `Fabius.thueMorseGammaTower_dyadic`. For $a > 0$ and $r \in \mathbb{N}_0$, the first two identify

$$\partial_s \tilde{D}(-r, a) = (-1)^r r! \text{ Mellin}(t \mapsto e^{-at} \mathcal{E}(t))(-r),$$

and the four identification theorems give its Mellin/integral form and the exact dyadic log and tower laws for every r . The definitions `Fabius.thueMorseGammaLog` and `Fabius.thueMorseGammaTower` are total for every real a , whereas their Mellin, integral, and dyadic theorems assume exactly $a > 0$. Formally, `Fabius.thueMorseGammaLog` is the chosen coordinate $L_r(a) := \partial_s \tilde{D}(-r, a)$, and the tower is defined by $T_r(a) := \exp L_r(a)$; no identity $L_r(a) = \text{Complex.log}(T_r(a))$ is proved. Thus the left side of (9.6) should be read as this chosen GammaLog coordinate in the Lean crosswalk, not as a branch theorem for the complex logarithm.

The current integrated tree adds the ninth public tower declaration `Fabius.ofReal_exp_mplimit_eq_gammaTower_div`: for $a, b > 0$, the complex embedding of $\exp(\text{mpLimit}(a, b))$ is $T_0(b)/T_0(a)$, exactly the formal ratio law for (9.10). The only unformalized assertions within [theorem 9.2](#) are now the parameter- a differential and iterated differential laws (9.8)–(9.9); they require differentiation under the Mellin integral.

Remark 9.4 (Why “Barnes-type”). The hierarchy is not asserted to be a classical Barnes multiple-gamma family. The name refers to the derivative ladder (9.8): higher levels integrate logarithms of lower levels, while the dyadic equation carries the automatic-sequence structure. This packaging appears to be absent from the inspected Atlas even though all analytic ingredients needed for its proof are present there.

10 Numerical verification and order diagnostics

All computations in this section used 100-decimal arithmetic. Product tails were extended far beyond the displayed precision. The calculations are reproducibility checks for the proved identities, not evidence on which the proofs depend.

10.1 Exact-identity residuals

Identity	level	argument	relative residual
oscillatory (3.1)	$m = 1$	$t = 0.37$	7.3×10^{-102}
oscillatory (3.1)	$m = 3$	$t = 2.125$	1.9×10^{-101}
oscillatory (3.1)	$m = 7$	$t = 5.75$	7.0×10^{-101}
oscillatory (3.1)	$m = 10$	$t = -3.4$	3.8×10^{-99}
Laplace (3.3)	$m = 4$	$s = 3.2$	5.6×10^{-101}
Laplace (3.3)	$m = 8$	$s = 11.5$	3.0×10^{-100}
Laplace (3.3)	$m = 12$	$s = 0.019$	9.4×10^{-96}

The residual level is consistent with accumulated arbitrary-precision rounding.

10.2 Observed convergence orders

A least-squares fit to the final five points of each curve gives:




Figure 1: Relative errors at $t = 5$. Every zeta correction or Richardson degree raises the power of 4^{-m} by one, exactly as predicted by [corollary 4.5](#) and [theorem 6.2](#).

method	fitted order in 4^{-m}	proved order
raw Q_m	1.00000013	1
zeta correction $N = 1$	2.00000006	2
zeta correction $N = 2$	3.00000007	3
zeta correction $N = 3$	4.00000008	4
q -Richardson $p = 1$	2.00000014	2
q -Richardson $p = 2$	3.00000013	3
q -Richardson $p = 3$	4.00000013	4

10.3 Weak expansion check

For $\varphi(y) = e^y$, all derivatives equal e^y , so [theorem 7.3](#) becomes a direct moment correction to the finite moment-generating function. The relative errors are:

m	raw spline	through μ_2	through μ_4	through μ_6
2	3.465×10^{-3}	4.569×10^{-6}	3.311×10^{-9}	1.540×10^{-12}
4	2.170×10^{-4}	1.789×10^{-8}	8.105×10^{-13}	2.357×10^{-17}
6	1.356×10^{-5}	6.991×10^{-11}	1.979×10^{-16}	3.597×10^{-22}
8	8.477×10^{-7}	2.731×10^{-13}	4.832×10^{-20}	5.489×10^{-27}
10	5.298×10^{-8}	1.067×10^{-15}	1.180×10^{-23}	8.375×10^{-32}

Each added even moment raises the order by one power of 4^{-m} .

10.4 Endpoint-tail coefficients

The absolute errors in approximating $\log G(-u)$ by the exact linear term plus the first N even corrections are:

u	$N = 0$	$N = 1$	$N = 2$	$N = 3$
0.02	5.56×10^{-6}	3.70×10^{-12}	5.60×10^{-18}	1.04×10^{-23}
0.2	5.56×10^{-4}	3.70×10^{-8}	5.60×10^{-12}	1.04×10^{-15}
1	1.39×10^{-2}	2.31×10^{-5}	8.71×10^{-8}	4.03×10^{-10}
3	1.23×10^{-1}	1.81×10^{-3}	6.12×10^{-5}	2.54×10^{-6}

The deterioration as u approaches the exact radius 4π is governed by the certified bound (5.7).

10.5 Lambert-depth experiment

For $x = 10^{-8}$,

$$\lambda(x) = 31.555232088424067978\dots, \quad s_0(x) = 3.1555232088424068 \times 10^9.$$




Figure 2: Tail error when $m = \lceil \lambda(x) + h \rceil$. Ceiling effects create the steps. The raw tail gains about one binary digit per extra factor; the centered tail gains about two, as proved in [theorem 8.1](#).

At $h = 8$, for example, $m = 40$ and $u \approx 2.8699 \times 10^{-3}$; the raw tail error is 1.43×10^{-3} while the centered tail error is 1.14×10^{-7} .

11 Consequences for the existing research program

11.1 A new exact-arithmetic/Fourier bridge

The Atlas treats the block polynomial P_m as a digital and algebraic object, while the Fourier-decay program treats dyadic sinc products through shell dynamics and lacunary sine products. Theorem 3.1 shows these are not merely related by analogy: P_m is the finite sinc product after one explicit normalization. This provides a direct translation dictionary:

finite digital statement	spectral statement
order- m zero of P_m at 1	removable t^{-m} normalization at the origin
Thue–Morse exponential sum	truncated up-function Fourier transform
Prouhet moment cancellation	high-order low-frequency cancellation
block length 2^m	product depth m
geometric error 4^{-m}	variance scale of omitted random tail

The shell-factorization results remain essential for large real t with m tied to the frequency shell. The present tail calculus addresses a different regime: once the first m factors have been isolated, it computes the remaining small-argument factor to arbitrary order. The two analyses are complementary rather than competing.

11.2 A new q -binomial role

The repository’s dyadic arithmetic naturally uses $q = 1/2$ because scales and evaluation points are dyadic. The present extrapolation naturally uses $q = 1/4$ because the centered error expansion is even. Gaussian binomials now appear in a second structural role:

$$\text{exact dyadic values: } q = \frac{1}{2}, \quad \text{spectral-tail acceleration: } q = \frac{1}{4}. \quad (11.1)$$

This distinction is conceptually useful. It prevents the q -binomial notation from being viewed as an incidental encoding of one exact formula; it reflects two different geometric lattices in the same theory.

11.3 A certified route from finite splines to up

The repeated-integration dossiers identify finite stages with compactly supported splines and prove convergence. The random-tail theorem strengthens the conceptual picture:

$$\text{limit law} = \text{finite spline law} + \text{scaled independent copy of the limit law}.$$

The weak expansion (7.7) then supplies every correction coefficient and a one-line remainder bound. For numerical integration against smooth observables, the result is immediately useful: one may integrate against the elementary finite spline u_m and add explicit derivative corrections, without ever sampling the limiting density directly.

11.4 A product engine for endpoint saddles

The endpoint theory’s difficult part is the periodically modulated saddle, not the mere existence of the product. Nevertheless, every saddle evaluation repeatedly calls the negative-Laplace product and its derivatives. Formula (5.8) suggests a hybrid implementation:

1. use the existing lower-Lambert phase and periodic saddle jets to choose s ;
2. choose m from the centered depth law (8.8);

3. evaluate the exact finite Thue–Morse block and a few Bernoulli corrections;
4. differentiate the finite factorization rather than the infinite product.

Because $u = s/2^m$ is small by construction, all differentiated correction series remain in a controlled analytic disk. This could turn the qualitative endpoint expansion into a certified high-precision algorithm with transparent complexity.

11.5 Formalization leverage

The main new identities have unusually short dependency chains. The finite bridges are finite-product algebra. Tail self-similarity is reindexing. The weak expansion is Taylor’s theorem plus independence and symmetry. The q -Richardson weights are finite Lagrange interpolation. Only the analytic logarithm and explicit zeta remainder require a substantial complex-analysis layer. This makes the package a plausible next formalization target — and in large part it is not a target at all but already done: as the *Formal version* paragraphs throughout [sections 2 to 4, 6 and 9](#) record, the block identity, both finite bridges in total normalized complex form, with denominator-free equalities at the origin and quotient corollaries away from zero, both halves of tail self-similarity, the Fourier–Laplace rotation, the Lagrange characterization of the q -Richardson weights with all three of its moment identities together with their closed q -binomial form (compiled in `GeometricQBinomialLagrange.lean`), every spectral level of the GammaLog/tower Mellin and dyadic laws together with the level-zero master-product ratio, and now the entire zeta–Lambert tail calculus — the Euler–zeta logarithm, the all-orders expansion, the certified remainder, and the exact prefix–corrections–tail split (`EulerLogTransform.lean`, `SincZetaSeries.lean`, `SincZetaDyadic.lean`, `SincZetaRemainder.lean`) — already have Lean counterparts, several of them more general than the statements printed here. The abstract weighted Richardson tail is now formal under unconditional summability in `AnalyticSeriesFilter.lean`, and `RvachevQBinomialFilter.lean` instantiates it exactly for the actual infinite Rvachev product. The identification with the finite prefixes Q_m used in [theorem 6.2](#), ordered and uniform remainder bounds, and the probabilistic layer are the remaining formalization targets.

12 Lean formalization roadmap

The finite-block calculus has been deliberately organized so that the algebraic core can be formalized before the analytic refinements. The following dependency table separates the short, high-confidence lemmas from the portions that require additions to the surrounding analysis library.

Several rows have since been overtaken by the library and are retained because the table is also a dependency map. The finite Thue–Morse product (2.1) is `Fabius.prod_one_sub_pow_eq_sum_thueMorseSign`, proved over an arbitrary commutative ring, not by the coefficient induction suggested below. The finite sinc and Laplace bridges (3.1)–(3.3) are now fully discharged in `ThueMorseComplexProductBridge.lean`. Their primary equalities are complex and total at the origin; the factors t^m and s^m expose the removable zeros, the prefix-at-zero simp laws give normalized value one, and the printed quotient normalizations are supplied as corollaries under the corresponding nonzero hypothesis. Exact tail self-similarity (3.4)–(3.5) is fully discharged, and the reusable reindexing lemma the row asks for is exactly what was proved: `Fabius.tprod_geom_scale_pow` is stated for an arbitrary function into a commutative topological monoid. Of the q -Richardson row, the interpolation characterization (6.5) and the moment identities (6.6)–(6.7), the forward Gaussian-binomial normalization (6.4), and the complete coefficientwise action on formal power series are now delivered at arbitrary ratio. The last item is the generic diagonal-filter API of `FinitePowerSeriesFilter.lean` and its geometric residual specialization `GeometricPowerSeriesFilter.lean`. The companion `AnalyticSeriesFilter.lean` evaluates an abstract coefficient series directly and gives the exact Gaussian tail whenever every contributing sample is unconditionally summable; zero

weights impose no unweighted convergence condition. It excludes conditionally-only boundary series. The separate `RvachevQBinomialFilter.lean` connects the named even modes to the actual infinite Rvachev product, but it does not identify the finite prefixes Q_m or their quotient and Bell coefficients. The Euler–zeta logarithm and certified zeta remainder have likewise since been delivered in `SincZetaSeries.lean` and `SincZetaRemainder.lean`, and the random-tail law exists in measure form in `MeasureRefinement.lean`. What remains includes the finite-prefix Richardson identification, weighted uniform remainder and convergence theorem, the random-variable weak expansion, the Lambert depth law, and the parameter- a differential and iterated differential Gamma-tower ladder. In particular, neither the formal filter identity nor the abstract unconditionally summable tail implies the analytic bounds or brackets stated later in this part.

Target	Minimal ingredients	Suggested proof strategy and principal risk
Finite Thue–Morse product (2.1)	finite products, binary recursion, geometric sums	Prove $P_{m+1}(z) = P_m(z)(1 - z^{2^m})$ and identify coefficients by induction. This should be almost entirely algebraic.
Finite sinc bridge (3.1)	complex exponential identities, finite-product reindexing	<i>Delivered</i> in <code>ThueMorseComplexProductBridge.lean</code> : the all-level total equality is <code>Fabius.thueMorseBlock_cexp_eq_sincPrefix</code> , its positive-level atlas wrapper is <code>Fabius.thueMorseBlock_cexp_eq_sincPrefix_of_pos</code> , and the quotient and origin laws are <code>Fabius.shiftedComplexSincPrefix_eq_thueMorseBlock_cexp_of_pos</code> and <code>Fabius.shiftedComplexSincPrefix_apply_zero</code> .
Finite Laplace bridge (3.3)	complex exponentials, finite products	<i>Delivered</i> in <code>ThueMorseComplexProductBridge.lean</code> by <code>Fabius.thueMorseBlock_exp_neg_eq_laplacePrefix</code> , with quotient corollary <code>Fabius.complexLaplacePrefix_eq_thueMorseBlock_exp_neg</code> and origin law <code>Fabius.complexLaplacePrefix_apply_zero</code> ; the same module proves the finite Fourier–Laplace rotation in <code>Fabius.complexLaplacePrefix_eq_exp_mul_shiftedComplexSincPrefix</code> .
Exact tail self-similarity (3.4)–(3.5)	convergence of infinite products, shifted indices	Establish a reusable infinite-product reindexing lemma. Existing repository convergence results should discharge the side conditions.
Euler–zeta logarithm (4.1)	local analytic logarithm, Euler’s sine product, normal convergence	<i>Delivered</i> in <code>SincZetaSeries.lean</code> (exponential and real-logarithm forms), on the general Euler log transform of <code>EulerLogTransform.lean</code> ; the all-orders identity (4.3) is <code>SincZetaDyadic.lean</code> .
Certified zeta remainder (4.6)	positivity, monotonicity of $\zeta(2r)$, geometric series	<i>Delivered</i> in <code>SincZetaRemainder.lean</code> : the verbatim bound, real-argument nonnegativity, and the exact prefix–corrections–tail split.

Target	Minimal ingredients	Suggested proof strategy and principal risk
q -Richardson weights (6.4)–(6.7)	finite Lagrange interpolation, q -Pochhammer algebra	<i>Delivered</i> in <code>GeometricQBinomialLagrange.lean</code> : zero-extended numerator in both forward and reversed Gaussian forms, the reversed denominator-free polynomial, normalized weight under node injectivity, and every positive residual moment. <code>GeometricPowerSeriesFilter.lean</code> further packages the entire coefficient-wise filter, while <code>AnalyticSeriesFilter.lean</code> supplies the exact unconditionally summable analytic Gaussian tail and its shifted-scale form, with convergence required only for nonzero-weight samples. <code>RvachevQBinomialFilter.lean</code> instantiates the actual infinite Rvachev product, but not the finite-prefix Q_m family in the displayed theorem. Conditional boundary summation, that finite-prefix coefficient bridge, and ordered or uniform remainder bounds are not part of this row. The case $q = 1/2$ is linked to the Toeplitz row.
Random-tail law (7.2)	independent random variables, almost-sure summation, equality in law	<i>Delivered</i> in measure form in <code>MeasureRefinement.lean</code> , by exactly the characteristic-function route this row predicted: the refinement equation and its m -fold iterate, via <code>charFun_conv</code> and the formal Fourier product.
Weak expansion (7.7)	Taylor theorem with remainder, expectation inequalities, symmetry	Condition on the finite sum; use bounded support to make the remainder estimate immediate.
Lambert depth law (8.7)–(8.8)	branch W_{-1} , ceilings, local exponential estimates	The analytic estimate is elementary after the definition of λ ; the likely missing piece is a convenient formal API for the lower real Lambert branch.
Thue–Morse gamma tower (9.6)–(9.9)	reciprocal-Gamma first jets, Atlas Dirichlet continuation, parameter differentiation	<i>Mellin and dyadic layers delivered</i> : <code>ReciprocalGammaJets.lean</code> proves the exact simple reciprocal-Gamma jet at every nonpositive integer, and <code>ThueMorseGammaTower.lean</code> proves every continuation first jet, the Mellin/integral <code>GammaLog</code> levels, both dyadic laws, and the level-zero master-product ratio. Only (9.8)–(9.9) remain; they need a parameter-under-the-Mellin-integral theorem.

A practical promotion sequence is therefore

$$\begin{array}{c} \text{finite bridges} \longrightarrow \text{tail reindexing} \longrightarrow \text{probabilistic splitting} \\ \longrightarrow q\text{-interpolation} \longrightarrow \text{analytic logarithms} \end{array}$$

The first four stages already yield exact identities and acceleration mechanisms. The final stage adds the sharp radius and explicit zeta remainder. This ordering minimizes the amount of new complex analysis that must be trusted at any one time. Every stage of the boxed chain now has a formal counterpart (see the *Formal version* paragraphs above): the finite bridges, tail reindexing, q -interpolation, and analytic logarithms directly, and the probabilistic splitting twice over – in measure form (`MeasureRefinement.lean`) and, through the affine bridge $\mu_{\text{up}} = \text{law}(X_\infty) \circ (2 \cdot - 1)^{-1}$, where $X_\infty = (Y_\infty + 1)/2$, as the equality-in-law splitting of the random series itself (`RandomSeriesLaw.lean`). Every stage of the chain is formal.

Formalization obligation 12.1 (Canonical normalization theorem – discharged). The repository uses several Fourier-transform conventions across historical manuscripts. A promoted theorem should package the convention explicitly and derive all rescalings between F , up , Q_∞ , and the characteristic function of Y_∞ . Doing so would prevent sign and 2π factors from leaking into later shell or saddle arguments.

This obligation is now discharged: `Fabius.normalization_dictionary` in `NormalizationDictionary.lean` is the promoted package. Its four conjuncts fix the convention $\widehat{\text{up}}_{2\pi}(z) = \int_{\mathbb{R}} \text{up}(t) e^{-2\pi i t z} dt$ and derive every rescaling simultaneously: the function-level fold $\text{up} = F(1 - |\cdot|)$ (also promoted separately as `Fabius.rvachevUp_eq_fabiusReal_one_sub_abs`), the product identity $\mathcal{F}_{2\pi}[\text{up}] = Q_{\infty}$ on all of $\mathbb{C}$, the centered characteristic function $\mathbb{E}[e^{itY_{\infty}}] = \widehat{\text{up}}_{2\pi}(t/(2\pi))$, and the uncentered $\mathbb{E}[e^{itX_{\infty}}] = e^{it/2} \widehat{\text{up}}_{2\pi}(t/(4\pi))$. The individual legs that fix the convention remain available: `Fabius.rvachevFourier` is $\int_{\mathbb{R}} \text{up}(t) e^{-2\pi i t z} dt$, `Fabius.rvachevFourier_eq_product` converts it into Q_{∞} , and `Fabius.complexGeneratingFunction_eq_fourier_analytic` together with `Fabius.complexGeneratingFunction_neg_eq_tprod` carries the rotation to G . The characteristic-function leg is now formal as well: `Fabius.charFun_weightedSumDistribution` evaluates $\mathbb{E}[e^{itX_{\infty}}] = e^{it/2} \widehat{\text{up}}_{2\pi}(t/(4\pi))$ for the uncentered $X_{\infty} = (Y_{\infty} + 1)/2$, while `Fabius.rvachevMeasure_charFun_pos` gives the centered identity $\mathbb{E}[e^{itY_{\infty}}] = \widehat{\text{up}}_{2\pi}(t/(2\pi))$. These fix the rescaling between Mathlib’s characteristic-function convention and `Fabius.rvachevFourier`, with the $F \leftrightarrow \text{up}$ passage carried at measure level by `Fabius.rvachevMeasure_eq_map_weightedSum` and its CDF identity `Fabius.rvachevMeasure_Iic`. The packaging the obligation asked for is exactly the four conjuncts of `Fabius.normalization_dictionary`: later shell or saddle arguments can cite the dictionary instead of reassembling sign and 2π conventions at each use.

Formalization obligation 12.2 (Reusable geometric-tail API – discharged). The identity “tail equals a scaled copy of the whole object” occurs simultaneously for products, convolutions, random series, and cumulant generating functions. A small abstract library for geometrically self-similar tails would make the dyadic case a specialization rather than four separate proofs.

The product quarter of this obligation is discharged. `GeometricScaleProducts.lean` is exactly the abstract library asked for on the product side: `Fabius.tprod_geom_scale` and `Fabius.tprod_geom_scale_pow` are proved for an arbitrary function from a field into a commutative topological monoid and an arbitrary scale, with the dyadic sinc shells obtained as specializations, exactly as the obligation proposes. The logarithmic face of the same abstraction is now `EulerLogTransform.lean`: the transform $\sum_i \log(1 - a_i) = -\sum_{r \geq 1} (\sum_i a_i^r)/r$ for an arbitrary absolutely summable family with every norm below one, over an arbitrary index type, with the geometric family – the self-similar tail par excellence – as its first instance. (The negative-Laplace tail (3.5) is formal too, but as an independent induction on the functional equation rather than as an instance of that API – which is itself a small illustration of the duplication the obligation is about.) The convolution quarter is now discharged in the same abstract style: `GeometricConvolutionTails.lean` builds the prefix system $P_0 = \delta_0$, $P_{m+1} = \nu * T_* P_m$ of an arbitrary digit measure under an arbitrary measurable additive endomorphism and proves, by pure convolution algebra, the cocycle law $P_{m+n} = P_m * (T^m)_* P_n$ (`Fabius.digitPrefix_add`) and the geometric tail law – from the one-step refinement $\mu = \nu * T_* \mu$ alone, $\mu = P_m * (T^m)_* \mu$ for every m (`Fabius.self_similar_conv_iterate`) – together with the real-scale specialization and the finite product formula $\text{charFun}(P_m)(t) = \prod_{k < m} \text{charFun } \nu(c^k t)$ (`Fabius.charFun_mulPrefix`). The dyadic instance (7.3) is a specialization: `MeasureRefinement.lean` now derives the measure-level random-tail law (7.2) from the refinement equation through this API (see the *Formal version* paragraph in section 7), and the random-variable representation underlying it is `ProbabilityRepresentation.lean`, with the equality-in-law splitting in `RandomSeriesLaw.lean`. The cumulant quarter is now formal as well: `GeometricCgfTails.lean` proves the moment tail law unconditionally (the moment generating function of a convolution factors with no integrability hypotheses, `Fabius.mgf_id_conv` and `Fabius.mgf_id_self_similar`), the cumulant tail law $K_{\mu}(t) = \sum_{k < m} K_{\nu}(c^k t) + K_{\mu}(c^m t)$ under finite exponential moments (`Fabius.cgf_id_self_similar`), and the dyadic instance with every hypothesis discharged by compact support (`Fabius.cgf_uniformDigitPrefix`, `Fabius.cgf_rvachevMeasure_self_similar`). With that, each of the four faces has the abstract library the obligation asked for, with the dyadic case a specialization, and the obligation is discharged.

The single statement instantiating to all four faces at once now exists as well: `GeometricTailDictionary.lean` proves `Fabius.geometric_tail_dictionary` — from the one-step refinement $\mu = \nu * (c \cdot)_* \mu$ alone, one theorem with four conjuncts: the measure face $\mu = P_m * (c^m \cdot)_* \mu$, the product face $\varphi_\mu(t) = (\prod_{k < m} \varphi_\nu(c^k t)) \varphi_\mu(c^m t)$ (whose assembly `Fabius.charFun_self_similar` it also provides), the unconditional moment face, and the cumulant face under finite exponential moments — with the dyadic instance `Fabius.geometric_tail_dictionary_up` carrying every hypothesis discharged by compact support. The operator-exponential display (7.12) is now formal at exactly the level the remark asserts it: `SmoothingOperatorExponential.lean` proves that for any zero-constant exponent series over a rational algebra, the smoothing operator of its generated exponential acts on every polynomial as the finite (degreewise-nilpotent) exponential of the exponent’s operator, $\mathcal{S}_{\exp K} q = \sum_{d \leq \deg q} \frac{1}{d!} \mathcal{S}_K^d q$ (`Fabius.seriesDerivOp_expSeries`, with the coefficient identity `Fabius.expCoeff_eq_sum_coeff_pow` and the vanishing of low coefficients of powers of a zero-constant series) — expanding the exponential changes cumulants into moments exactly as claimed, with the `Fabius` instance available through the moment–cumulant dictionary of `MomentCumulantAlgebra.lean`. Only the distributional reading of (7.12) against the spline densities remains analytic.

13 Frontier conjectures and next deductions

The proved results isolate several questions which are both sharper and more tractable than the broad open-ended prompts in the older frontier notebooks. The statements below are intentionally separated from the theorem package.

Conjecture 13.1 (Certified endpoint evaluator with near-optimal product depth). *There is a fully explicit algorithm which, for sufficiently small $x > 0$ and target relative error η , evaluates $F(x)$ using*

$$-\frac{W_{-1}(-x \log 2)}{\log 2} + \frac{1}{2} \log_2 \frac{1}{\eta} + \mathcal{O}(1) \quad (13.1)$$

product layers, together with a bounded-order periodic saddle jet, and whose final error is rigorously at most η . The coefficient $1/2$ in the precision surcharge is optimal among methods which first center the omitted product but do not cancel its quadratic cumulant.

The depth estimate itself is proved in [theorem 8.1](#). What remains is to propagate the product error, its differentiated versions, and the periodic saddle displacement through the complete inversion formula. The conjectured optimality should follow from the nonzero variance $\kappa_2 = 1/9$.

Conjecture 13.2 (Cumulant-cancelled endpoint hierarchy). *After removing the first N even cumulants of the centered tail in addition to the linear factor, the extra depth required for tolerance η is*

$$\frac{1}{2N+2} \log_2 \frac{1}{\eta} + \mathcal{O}_N(1). \quad (13.2)$$

Equivalently, the N -corrected finite block should require only one additional product layer for every $2N+2$ requested binary digits of tail accuracy.

This is already a theorem for the isolated product, by (5.7); the conjectural part is that the same saving survives uniformly through the full saddle inversion, including all needed derivatives and the periodic phase.

Conjecture 13.3 (Strong spline correction away from knots). *Fix $N \geq 0$ and a compact interval $K \subset (-1, 1)$. After excluding a neighborhood of the knots of the finite spline u_m whose width is comparable to 2^{-m} , the distributional expansion (7.10) admits a pointwise or Hölder-norm version through order N , with remainder $\mathcal{O}_K(4^{-(N+1)m})$ after the natural derivative scaling.*

A global pointwise statement cannot be obtained by naively differentiating u_m : finite splines have knots, whereas up is smooth. The random-tail convolution indicates two viable repairs. One may either retain the exact final mollifier $2^m \text{up}(2^m \cdot)$, or state the expansion on cells separated from their endpoints. Establishing the best norm and boundary layer is a concrete approximation-theoretic problem.

Conjecture 13.4 (Dual- q transmutation). *There is a direct coefficient-level transform taking a substantial class of the repository's $q = 1/2$ exact dyadic formulae into $q = 1/4$ spectral-tail acceleration identities, with the squaring $q \mapsto q^2$ explained by parity or centering rather than by an ad hoc substitution.*

Evidence is structural: the exact dyadic theory resolves single binary scales, while every centered Fourier or probability correction is even and therefore sees squared scales. A useful form of the conjecture would turn Gaussian-binomial coefficient extraction at $q = 1/2$ into geometric-grid interpolation at $q = 1/4$.

Conjecture 13.5 (Automatic Barnes hierarchy). *The functions Γ_r^{TM} possess canonical Weierstrass products and asymptotic normalizations for which the dyadic law (9.7) and differential ladder (9.8) uniquely characterize the hierarchy, analogously to the role of convexity or normalization in classical gamma-type characterizations.*

The Mellin formula gives positivity and complete-monotonicity information for suitable signed derivatives, while the dyadic equation supplies an automatic analogue of a multiplication formula. The remaining difficulty is to choose a normalization that removes the freedom to multiply by functions invisible to the differential recursion.

13.1 General radix and nonuniform tails

The base- b theorem (4.11) shows that the zeta–Lambert mechanism does not depend on binary arithmetic. What is special to base two is the simultaneous existence of the Thue–Morse block polynomial, the simple two-branch functional equation, and the Fabius/up probability model. A next generalization is to replace the uniform summands by centered compactly supported variables X_k at scales b^{-k} . If their common characteristic function has a nonzero analytic germ at the origin, then

$$\log \frac{\Phi_\infty(t)}{\Phi_m(t)} = \sum_{r \geq 1} c_r t^r \frac{b^{-r(m+1)}}{1 - b^{-r}}.$$

Thus Lambert kernels are universal for geometric random tails; the zeta values in the present report are the special signature of the uniform law through Euler's sine product. The classification of digital substitutions whose finite block polynomial exactly realizes such a characteristic product is a natural interface between automatic sequences and compactly supported refinable laws.

14 Conclusion

The documentation corpus already supplied the large pieces: a careful separation of the Fabius and up-function normalizations, an exact dyadic arithmetic theory rich in q -binomial coefficients, a probabilistic convolution model, a detailed analysis of the infinite sinc product, corrected lower-Lambert endpoint asymptotics, and a broad Thue–Morse formula atlas. The missing junction was finite-block calculus.

One polynomial now connects the principal domains:

$P_m(z) = \sum_{\substack{n < 2^m \\ }} \varepsilon_n z^n$	finite digital block finite Fourier sinc product finite negative-Laplace product
--	--

after explicit phase and scale normalizations. The omitted object is always a scaled copy of the full object. That observation produces exact all-order tails, zeta–Lambert coefficients, Gaussian-binomial extrapolation, random-tail deconvolution, and Lambert- W complexity laws from a common source rather than as isolated formulas.

The principal proved advances are:

1. exact oscillatory and decaying substitutions converting a length- 2^m Thue–Morse block into the first m factors of the Fourier and endpoint products;
2. a convergent all-orders logarithmic tail with exact radius and an explicit geometric remainder bound;
3. arbitrary-order corrected finite-block approximants and a separate q -Richardson scheme whose weights are Gaussian binomials at $q = 1/4$;
4. an exact independent-copy decomposition of the omitted random tail, yielding a certified weak differential expansion of the up-law in finite uniform splines;
5. raw, centered, and higher-correction product-depth laws on the canonical lower-Lambert saddle scale;
6. a Thue–Morse gamma tower obtained by differentiating the entire shifted Dirichlet function at its trivial zeros.

These results do not replace the repository’s sharp shell dynamics or periodically modulated endpoint saddle. They add a complementary local engine: after a finite block or finite shell has been isolated, the remainder is computable to arbitrary order with transparent constants. That engine is algebraically short, numerically stable after centering, and unusually suitable for formal verification.

A Coefficient and extrapolation tables

A.1 Zeta–Lambert and centered Laplace coefficients

Write

$$a_r = \frac{\zeta(2r)}{r\pi^{2r}(4^r - 1)}, \quad c_r = \frac{B_{2r}}{2r(2r)!(4^r - 1)}.$$

Then

$$\log \frac{Q_\infty(t)}{Q_m(t)} = - \sum_{r \geq 1} a_r t^{2r} 4^{-rm}, \quad \log C(-u) = \sum_{r \geq 1} c_r u^{2r}.$$

The first eight exact coefficients are:

r	a_r	c_r
1	1/18	1/72
2	1/2700	−1/43200
3	1/178605	1/11430720
4	1/9639000	−1/2467584000
5	1/478533825	1/490018636800
6	691/15688261338750	−691/64259118443520000
7	2/2092151286225	1/17138903336755200
8	3617/170727360353550000	−3617/11188788288130252800000

The numerators 691 and 3617 are the familiar Bernoulli numerators; here they enter through the uniform-tail cumulants rather than through an independent number-theoretic construction.

A.2 Analytic error germ

For

$$H_t(r) = \frac{1}{Q_\infty(t\sqrt{r})} = \sum_{n \geq 0} h_n(t) r^n,$$

the first coefficients are monomials in t^2 :

$$\begin{aligned} h_0(t) &= 1, \\ h_1(t) &= \frac{t^2}{18}, \\ h_2(t) &= \frac{31t^4}{16200}, \\ h_3(t) &= \frac{2347t^6}{42865200}, \\ h_4(t) &= \frac{1904369t^8}{1311675120000}, \\ h_5(t) &= \frac{3309760579t^{10}}{88561680752160000}. \end{aligned}$$

They are complete exponential Bell polynomials in $a_1 t^2, a_2 t^4, \dots$. Consequently, all leading constants in (6.10) are exact rational multiples of $Q_\infty(t) t^{2p+2}$.

A.3 Gaussian-binomial extrapolation weights

At $q = 1/4$, the first six rows are:

$$p = 0 : \quad (1),$$

$$\begin{aligned}
p = 1 : & \quad \left(-\frac{1}{3}, \frac{4}{3} \right), \\
p = 2 : & \quad \left(\frac{1}{45}, -\frac{4}{9}, \frac{64}{45} \right), \\
p = 3 : & \quad \left(-\frac{1}{2835}, \frac{4}{135}, -\frac{64}{135}, \frac{4096}{2835} \right), \\
p = 4 : & \quad \left(\frac{1}{722925}, -\frac{4}{8505}, \frac{64}{2025}, -\frac{4096}{8505}, \frac{1048576}{722925} \right), \\
p = 5 : & \quad \left(-\frac{1}{739552275}, \frac{4}{2168775}, -\frac{64}{127575}, \frac{4096}{127575}, -\frac{1048576}{2168775}, \frac{1073741824}{739552275} \right).
\end{aligned}$$

The alternating signs are unavoidable because the formula extrapolates from positive nodes $q^m, \dots, q^{m+p}$ to the endpoint 0. For numerical work at very high p , barycentric or compensated summation should be used rather than forming the rational coefficients in floating point.

B Reproducibility protocol

The numerical checks used direct high-precision arithmetic and did not fit coefficients that were subsequently inserted into proofs. The following pseudocode records the independent verification path.

Listing 1: Independent verification skeleton

```

set_precision(100)

def epsilon(n):
    return (-1) ** binary_digit_sum(n)

def P(m, z):
    return product(1 - z ** (2 ** j) for j in range(m))

def Q_finite(m, t):
    return product(sinc(t / (2 ** k)) for k in range(1, m + 1))

def Q_infinite(t, cutoff):
    return product(sinc(t / (2 ** k)) for k in range(1, cutoff + 1))

# Exact finite identities
for (m, t) in test_grid:
    lhs = P(m, exp(1j * t / (2 ** (m - 1))))
    rhs = ((-1j) ** m * t ** m * 2 ** (-m * (m - 1) / 2)
           * exp(1j * t * (1 - 2 ** (-m))) * Q_finite(m, t))
    report(relative_residual(lhs, rhs))

# All-order tail and q-Richardson slopes
for m in depth_grid:
    exact = Q_infinite(t, large_cutoff)
    for N in correction_orders:
        approx = Q_finite(m, t) * exp(-sum(a(r) * t ** (2*r)
                                           * 4 ** (-r*m)
                                           for r in range(1, N + 1)))
        save_error(m, N, relative_error(approx, exact))
    for p in extrapolation_orders:
        approx = sum(weight(p, j) * Q_finite(m + j, t)
                     for j in range(p + 1))

```

```

    save_error(m, p, relative_error(approx, exact))
fit_final_log_slopes()

# Random-tail weak expansion,  $\phi(y)=\exp(y)$ 
for m in depth_grid:
    finite_mgf = product(sinh(2 ** (-k)) / (2 ** (-k))
                        for k in range(1, m + 1))
    corrected = finite_mgf * sum(moment(2*j) * 4 ** (-j*m)
                                / factorial(2*j) for j in range(N + 1))
    compare_with_infinite_mgf(corrected)

```

The oscillatory and Laplace finite identities were tested at levels up to $m = 12$ with relative residuals between 10^{-96} and 10^{-102} . Regression was used only to diagnose the convergence exponent after the formulas had been proved; the fitted base-4 orders were 1, 2, 3, 4 to roughly seven significant digits for the raw, first-, second-, and third-order schemes. The figures embedded in this report were generated from the same data.

C Complete corpus ledger and deduplication map

C.1 Documentation treated as authoritative at the audit snapshot

The audit covered every TeX manuscript in the living documentation tree. At the snapshot date, the principal paths and roles were:

Path or group	Role in the audit
Fabius_Function_and_Rvachev_Up /Fabius_Function_and_Rvachev_Up. tex	Primary synthesis: normalization, probability, Fourier theory, exact dyadic values, and corrected endpoint asymptotics.
FabiusFunction_Mathematical_Gl ossary/FabiusFunction_Mathemat ical_Glossary.tex	Terminology and notation cross-check.
Non_Elementarity_of_the_Fabius _Function/Non_Elementarity_of_ the_Fabius_Function.tex	Differential-algebraic and non-elementarity context.
fabius_lean_walkthrough/fabius _lean_walkthrough.tex	Map between human mathematics, Lean files, dependencies, and proof status.
semi-formalized-research-front iers/semi-formalized-research-f rontiers.tex	Consolidated gap register and non-formalized frontier synthesis.
semi-formalized-research-front iers/drafts/thue-morse/Thue_Mo rse_Atlas_and_Frontiers/Thue_M orse_Atlas_and_Frontiers.tex	Current consolidated Thue–Morse volume; Part I is the comprehensive formula atlas and principal source for the block polynomial and shifted Dirichlet continuation.

Path or group	Role in the audit
semi-formalized-research-frontiers/drafts/rvachev_up_fourier_decay/	Twelve then-active manuscripts: the initial wave, comparative and second-wave audits, a gentle guide, the asymptotic-product problem statement, successive waves 2–8, and the final resolution. They were compared rather than collapsed in that audit because their hypotheses and correction histories differ. They are now consolidated in rvachev_up_fourier_decay/Rvachev_Up_Fourier_Decay/, with exact historical bytes preserved at immutable commit 2e3567feb149.
papers/AIF_76_21-38/	Original and corrected TeX transcriptions of Arias de Reyna’s compact-support paper.
papers/arXiv-1502.06738/	Source TeX for Aistleitner–Hofer–Larcher on parametric Thue–Morse sequences and lacunary trigonometric products.
papers/arXiv-1702.05442/	Source TeX for Arias de Reyna’s English translation of the 1982 compact-support article.
papers/arXiv-1702.06487/	Source TeX for Arias de Reyna’s arithmetic study of dyadic Fabius values.
papers/arXiv-2211.13570/	Source TeX for Tóth’s shifted and combined Thue–Morse Dirichlet series.
papers/ubiq15/ubiq15-corrected.tex	Corrected source for imported automatic-sequence background.

At the immutable pre-consolidation snapshot 2e3567feb149, the inspected Fourier-decay sub-directories were Rvachev_Up_Fourier_Decay-2, Rvachev_Up_Fourier_Decay-3, Comparative_Audit, Gentle_Guide, Second_Wave_Audit, asymptotic-decay-rate-of-an-infinite-product-of-sinc-functions, rvachev_up_fourier_decay-1, and rvachev_up_fourier_decay-4 through rvachev_up_fourier_decay-8. Each contained one principal TeX manuscript. Those live copies are now retired; current navigation uses rvachev_up_fourier_decay/Rvachev_Up_Fourier_Decay/Rvachev_Up_Fourier_Decay.tex.

C.2 Frozen archive represented through provenance

The archive README states that entries preserve the first-introduced historical document, that later edits are not separately archived, and that byte-identical copies already present in the living tree were pruned. Accordingly, the audit used the following unique historical entries as a provenance map rather than counting every renamed or duplicated copy as new mathematical evidence.

Archive category	Frozen entries
Early notes	FabiusAsymptotic; K-foldsummationoverthesignedThue–Morse sequence.
Standalone studies	Non_Elementarity_of_the_Fabius_Function; Small_Argument_Asymptotics; fabius-inverse-dyadic-closed-form.
Primary exposition	Fabius_Function_and_Rvachev_Up; Fabius_Function_and_Rvachev_Up-2; Fabius_Function_and_Rvachev_Up-3; Fabius_and_Rvachev_Up-2; Fabius_and_Rvachev_up; Fabius_function_and_up_function.

Archive category	Frozen entries
Primary supplements	Fabius_Function_Missing_Parts; Fabius_Function_Missing_Parts-2.
q -calculus	Fabius_q_Special_Functions; fabius-q-special-functions.
Lean walkthrough	Two frozen walkthrough variants recorded under the archive’s lean-walkthrough category.
Glossary	Two frozen glossary variants recorded under the archive’s glossary category.
Research frontiers	Fabius_Dyadic_Asymptotic_Bridge; Fabius_Dyadic_Formulae_and_Alternative_Representations; Fabius_Dyadic_Formulae_to_Asymptotics; Fabius_Dyadic_q_Connections; Fabius_Thue_Morse_Convergence_Rate; Fabius_Thue_Morse_Convergence_Rates; Primary_Exposition_Gap_Register; Repeated_Integration_and_Rvachev_Up; Rvachev_Up_from_Repeated_Integration; semi-formalized-research-frontiers; Thue_Morse_Formula_Atlas-1; Thue_Morse_Formula_Atlas-2; Thue_Morse_Formula_Atlas-3.

This ledger explains the phrase “all TeX documents” operationally: every living manuscript was inspected, every active mutually auditing draft was retained, every imported TeX paper was included, and frozen historical variants were read through their explicit provenance and content map so that superseded copies did not distort novelty checks.

References

- [1] V. Reshetnikov, *The Fabius Function and Rvachev’s Up-Function: Exact arithmetic, probability, Fourier analysis, and corrected sharp asymptotics*, living repository manuscript, 2026. https://github.com/VladimirReshetnikov/ProveIt/blob/main/Analysis/FabiusFunction/docs/Fabius_Function_and_Rvachev_Up/Fabius_Function_and_Rvachev_Up.tex
- [2] V. Reshetnikov, *Non-Formalized Research Frontiers for the Fabius Function*, living repository manuscript, 2026. <https://github.com/VladimirReshetnikov/ProveIt/blob/main/Analysis/FabiusFunction/docs/semi-formalized-research-frontiers/semi-formalized-research-frontiers.tex>
- [3] V. Reshetnikov, *The Thue–Morse Sequence: Formula Atlas and Fabius–Rvachev Frontier Results*, living repository manuscript, 2026; Part I contains the unified formula atlas. https://github.com/VladimirReshetnikov/ProveIt/blob/main/Analysis/FabiusFunction/docs/semi-formalized-research-frontiers/drafts/thue-morse/Thue_Morse_Atlas_and_Frontiers/Thue_Morse_Atlas_and_Frontiers.tex
- [4] V. Reshetnikov, *The Fourier Decay of Rvachev’s Up-Function: Final Resolution*, living repository manuscript, 2026. https://github.com/VladimirReshetnikov/ProveIt/blob/main/Analysis/FabiusFunction/docs/semi-formalized-research-frontiers/drafts/rvachev_up_fourier_decay/Rvachev_Up_Fourier_Decay/Rvachev_Up_Fourier_Decay.tex
- [5] V. Reshetnikov, *Archive of historical document versions*, provenance README, 2026. <https://github.com/VladimirReshetnikov/ProveIt/tree/main/Analysis/FabiusFunction/docs/archive>

- [6] J. Arias de Reyna, *An infinitely differentiable function with compact support: Definition and properties*, English translation of *Definición y estudio de una función indefinidamente diferenciable de soporte compacto*, Rev. Real Acad. Ciencias Madrid **76** (1982), 21–38; arXiv:1702.05442 (2017). <https://arxiv.org/abs/1702.05442>
- [7] J. Arias de Reyna, *Arithmetic of the Fabius function*, arXiv:1702.06487v3 (2017). <https://arxiv.org/abs/1702.06487>
- [8] C. Aistleitner, R. Hofer, and G. Larcher, *On parametric Thue–Morse sequences and lacunary trigonometric products*, arXiv:1502.06738v2 (2015). <https://arxiv.org/abs/1502.06738>
- [9] L. Tóth, *Linear Combinations of Dirichlet Series Associated with the Thue–Morse Sequence*, INTEGERS **22** (2022), Article 98; arXiv:2211.13570. <https://arxiv.org/abs/2211.13570>
- [10] R. M. Corless, G. H. Gonnet, D. E. G. Hare, D. J. Jeffrey, and D. E. Knuth, *On the Lambert W function*, Advances in Computational Mathematics **5** (1996), 329–359. <https://doi.org/10.1007/BF02124750>
- [11] J.-P. Allouche and J. Shallit, *Automatic Sequences: Theory, Applications, Generalizations*, Cambridge University Press, 2003.