

Local quantitative refinements of Gowers’s proof of Szemerédi’s theorem

Density transfer and phase-flat partitions; energies, restriction and dependent random choice;
cube and progression counts; floor patterns and the Proposition 17.7 phase extraction

Sixty-nine research manuscripts prepared for Vladimir Reshetnikov,
merged into one ProveIt research report (sources 01–69)

Sources: 6–7 October 2026 merged: 6–7 October 2026

Abstract

Gowers’s 2001 proof of Szemerédi’s theorem passes through many local estimates: phase-flat partitions and density increments (Section 2 and Lemma 5.15 of the paper), the energy and restriction steps of the inverse theorem (Proposition 6.1, Proposition 7.3, Lemma 7.4, Lemmas 9.2–9.4, 12.4, 14.7–14.8), cube and progression counts (Corollaries 3.3–3.6, Lemma 3.10) and floor-pattern counts (Corollaries 17.5–17.6, Proposition 17.7). This report merges five independent research manuscripts, all dated 6 October 2026 and prepared for Vladimir Reshetnikov, that sharpen these estimates against the corrected statements of ProveIt’s transcription and Lean development. The material is arranged in five thematic Parts. *Part I*: a sharp one-sided density–mass bound, proved in four sources; a sharp one-cell phase envelope; the lower constant $1/3$ (optimal) in Lemma 2.3; Corollary 2.4 with 36π and $3\alpha/4$; three density-sensitive Fourier increments, of which the increment $3\delta(1-\delta)\kappa/(4-(4-3\delta)\kappa)$, $\kappa = \eta/(2\delta(1-\delta))$, of one source is at least $3/2$ times those of the other two at their stated parameters and lengths of the same order (proved in the write); simultaneous phase-flat partitions with the optimal exponent $1/(m+1)$. *Part II*: a mixed, weighted extension of Proposition 6.1 on every finite abelian group with coefficient one; a cubic dependent-random-choice certificate giving $7\delta^3 n/(3\sqrt{6})$ in Lemma 7.4; an elementary Balog–Szemerédi–Gowers bound $(c/8, 2^{19}c^{-9})$; exact moment interpolation and a polynomial restriction theorem over prime fields. *Part III*: the exact fourth-order term of the centered cube count on every finite abelian group; the sharp gap $\|f\|_{U^3}^8 \geq 2\|f\|_{U^2}^8$ for real centered functions on odd-order groups, hence the sharp leading constant $6\sqrt{2}$ and, in the write, the answer $c_3^{\text{odd}} = 2^{-1/2}$ to one source’s question; the fractional correction exponent $16/3$; density-weighted progression counting. *Part IV*: the floor-pattern count C_k (two sources) with $|\mathcal{F}(\mathbf{r})| = (\prod_i 2r_i)C_k$, $2^{k(k-1)/2} \leq C_k$, $\log_2 C_k \sim k^2$ (using the classical threshold-function count), certified $C_1, \dots, C_4 = 1, 2, 10, 154$, and a local phase-extraction coefficient $2^{-k^2 - k \log_2 k - O(k)}$ at Proposition 17.7. *Part V*: the crosswalk to the formal ledger, the sources’ formalization plans, and their 46 research questions, deduplicated with credits. No source claims a new global bound for $r_k(N)$. The report is unrefereed. Of its statements only one, a lifting lemma from modular to integer progressions, is formalized, and it was formalized before the sources arrived; the report gains no formal status.

[*Second write, 6 October 2026: sources 06–10.*] Five further manuscripts of the same date are added to the same Parts, so the report now merges ten. *Part I* gains the exact optimum for the largest dense cell with a fixed number of cells (strictly better than the Lemma 5.15 replacement above for $\alpha < 1$), the optimal density cap $\delta + \delta S/(2\delta - S)$ and a density–length frontier; the sharp one-sided bound is now proved in six sources. *Part II* gains a chapter on Lemma 10.9: a transport-sensitive rigidity theorem with the sharp threshold $\delta < 1/(2k)$, which upgrades the

Freiman order from two to five under exactly the hypotheses of the Lean statement and to nine under further hypotheses of the construction; and an energy extraction with $(c/4, 2^{10}c^{-5})$ and a Freiman extraction giving $2^{-209}\gamma^{103}$ in place of $2^{-1882}\gamma^{1164}$ in Corollary 7.6. *Part III* gains the complete classification of the cube count through order five on every finite abelian group, with Stirling multiplicities and order-two and order-three torsion terms, a determinant-five configuration at order six, the refutation of the quartic lower bound with coefficient 12 (recorded with an explicit counterexample), a quadratic progression discrepancy, torsion-sensitive counting, and an exact dimension threshold for polynomial colourings. *Part IV* gains a second certificate of $C_4 = 154$ and carry entropy uniform in the modulus. *Part V* gains iteration potentials and 50 further questions (96 in all). No statement of the new sources is formalized.

[*Third write, 6 October 2026: sources 11–14.*] Four further manuscripts of the same date are added to the same Parts, so the report now merges fourteen. *Part I* gains complete simultaneous partitions with prescribed lengths $v - 1$ or v and exponent $1/(3q + 2)$ for q linear phases (no size threshold), random-slope obstructions with independent tolerances, and a joint phase-density-exception frontier; the sharp one-sided bound is now proved in seven sources, and the new Fourier increment is shown at the write to be dominated by an earlier one. *Part II* gains two chapters: sharp affine rigidity near the exact endpoint of Lemma 6.2 (cubics $4t - (10 + 2\kappa)t^2 + (6 + 2\kappa)t^3$ with coset extremizers, and the exact largest non-affine graph energy for $|G| \geq 34$, whose cutoff cannot simply be dropped: Remark W2), and sharp Freiman extension with polynomial Bohr radii in Section 10 ($3\varepsilon + 2\kappa < 1$, sharp; radius $\alpha^{19}/(3\pi 2^{149})$); and near-maximal energy forcing a unique nearby coset with error $(1 - \sqrt{1 - 4\varepsilon})/2$, which, composed at the write with the affine-rigidity bounds, answers one source’s stability question in substance (Remark W3). *Part III* gains the exact secondary coefficient $(9/4)^{1/3}$ of the three-cube extremal problem, which *settles* source 04’s Question 1 (Remark W4), the exponent six on groups of exponent five, centered three-cube orbit moments, interval discrepancy with the sharp exponent $1/4$, and exact supersaturation for polynomial colourings. *Part V* gains 45 further questions (141 in all). No statement of the new sources is formalized.

[*Fourth write, 6 October 2026: sources 15–20.*] Six further manuscripts of the same date are added to the same Parts, so the report now merges twenty. *Part I* gains partitions whose own steps control affine frequencies (every balanced exponent below $1/((q + 1)(42q + 1))$, conditionally on a cited recurrence theorem; an obstruction at $1/(3q + 1)$) and, in two independent proofs conditional on Maynard’s theorem, simultaneous polynomial partitions with exponents polynomial in the number of phases, which answer one source’s question, with matching-order obstructions. *Part II* gains three chapters: collision-sensitive amplification at Lemmas 14.6–14.8 ($\beta^{22 \cdot 4^k} \gamma^{44k4^k + 56 \cdot 2^k}$ for eight-arrangements, against $\beta^{28 \cdot 4^k} \gamma^{84k4^k}$ in the formal statement); polynomial-loss restriction at Section 9 ($\alpha^{16}\eta^{15}/(2 \cdot 18^{16})$ at order eight, which proves the prime-modulus case of the formal Lemma 9.3 with an explicit threshold); and the sharp degenerate-arrangement density $(k + C_d)/p$ at Lemma 15.4 with exact common sampling and shared-slice interpolation at Lemma 16.10 (sampling constant $\varphi^3 e^{-\varphi}$). *Part III* gains the cube expansion through order five twice more (independent derivations of an earlier source’s theorem) and quadratic progression counting with a coefficient strictly below an earlier source’s (proved at the write), whose exponent is sharp for indicator sets. *Part IV* gains Proposition 17.7 without carry patterns, with retention $(4/(3(k + 2)))^{k+2}$: at $k = 8$ it loses 29 bits, against 1,458 in the formal statement and 91 for the best earlier source in the same normalization (re-derived at the write). *Part V* gains 74 further questions (215 in all). No statement of the new sources is formalized.

[*Fifth write, 6 October 2026: sources 21–27.*] Seven further manuscripts of the same date are added to the same Parts, so the report now merges twenty-seven. *Part I* gains sharp thresholds for simultaneous linear-phase partitions with prescribed cell lengths: the numerical-semigroup conductor enters with leading constant one for lengths $L, L + 1$, and a mixed-radix obstruction, valid even with cell-dependent steps, *answers* one source’s question: the prescribed-length frontier $(q + 2)a + qb = 1$, hence the exponent $1/(3q + 2)$ of that source, is optimal. It also gains a chapter on Corollary 5.8 without a scale threshold: at degree one, gain $\alpha/3$ at length $(N/M)^{1/3}/8$ over every cyclic group, against $\alpha/16$ and $(N/M)^{1/16}/8$ in the open formal statement, whose degree-one instance holds trivially when $N/M \leq 2^{48}$ (proved at the write). *Part II* gains a

Section 7 chapter, mostly second routes to two earlier sources, with one new extraction for codomains $\mathbb{F}_p^d$ that does not depend on d ; three more treatments of degenerate arrangements at Lemma 15.4 (a degree-stratified finite bound; the leading coefficient $(k^2 + 3k)/p$ for $d = 1$; the codimension-one classification and a selection theorem for every d ; the exact second coefficient $(k^2 + 21k - 6)/(2q^2)$ for $d = 2$, with an exact eight-vertex count), where the leading coefficient $k + C_d$ is an earlier source's, re-derived twice independently; and a common base partition at Lemma 16.10 that keeps width $(m/3)^A$. *Part III* gains the sharp exponent-five coefficient $2^{3/4}/3$ with a uniform $O(u^8)$ error, in two independent proofs, which *answers* one source's question, the exact next coefficient $20/9$ on $\mathbb{Z}/5\mathbb{Z}$, and degree filtrations of polynomial colourings with an exact pointwise pattern minimum. *Part V* gains 74 further questions (289 in all). No statement of the new sources is formalized.

[*Sixth write, 6 October 2026: sources 28–33.*] Six further manuscripts of the same date are added to the same Parts, so the report now merges thirty-three. *Part I*: Corollary 5.8 at degree two without a scale threshold (gain $\alpha/3$ at length $(N/M)^{1/39}/17$ over every cyclic group), which with the degree-one chapter gives written proofs of the instances $k \leq 2$ of the open formal statement.

Part II: the exact optimal near-extremal U^3 profile $(1 - \sqrt{(1 + \sqrt{1 - 4\delta})/2})/2$ on every finite abelian group and sharp selected-graph stability at the Proposition 6.1 endpoint, with a second, cruder route; weights carried through the Section 7 restriction and a correlation-preserving exact restriction, which answers a question of an earlier source in a prime-cyclic form; a fixed-field restriction keeping every degenerate tuple with a threshold linear in $(\alpha\eta)^{-1}$; at Section 10 a logarithmic selection retaining $\alpha p(\alpha)/20$ ($p(\alpha)$ of order $\log^{-2}(1/\alpha)$) with $O(\alpha^{-3})$ frequencies and radius $\Omega(\alpha^5)$, and a fixed-radius translation estimate with a proof that exponential dependence on the rank is necessary; at Lemma 12.3 the linear branch $4\theta - 3$ and a sharp high-agreement repair of respected arrangements on full products. *Part III*: $\|f\|_{U^4}^{16} \geq (35/8) \|f\|_{U^2}^{16}$ for real centered f on odd-order groups and $c_d^{\text{odd}} \geq 1/3$ for every d ; with an exact two-harmonic witness and a certificate of the intake, $0.5587725 \leq c_4^{\text{odd}} \leq 0.6914416$ (later placed sources record better upper bounds), and a single cosine is not extremal for $d \geq 4$. *Part V* gains 69 further questions (358 in all). No statement of the new sources is formalized.

[*Seventh write, 6 October 2026: sources 34–39.*] Six further manuscripts of the same date are added to the same Parts, so the report now merges thirty-nine. *Part I*: the exact first-window missing-arc minimax $\rho_R(a) = 1/(R - \cot(a/2) \tan((R+1)a/2))$ and a replacement for Lemma 5.2 with frequency at most $\lfloor N/M \rfloor - 1$ and coefficient $|A|M/(N - M)$, which implies the formal statement, and a recurrence exponent $(1 - 1/(4k))/(2^{k-1} + 1)$ from the formal Weyl input; explicit polynomial phases forcing $D \mid h$ on every flat progression, with exact thresholds $DC(\mathcal{S})$ and the necessary frontier $(W + 2)a + \sum_i d_i b_i \leq 1$ for prescribed lengths. *Part II*: degree-weighted Fejér selection for Lemma 15.5 with field-size exponent $E_{k,m} \sim 2s/k$; the sharp odd-order U^3 stability modulus $\varepsilon/4 + 3\varepsilon^2/16 + (\sqrt{2}/64)\varepsilon^{5/2}$; a shared-anchor sampler for Lemma 16.10 that keeps sampled targets, with a closed solution of its finite minimax; the sharp nonuniform transport threshold $1/(2 + (2k - 2)\kappa)$ for every κ . *Part III*: $\|f\|_{U^2}^4 \leq (2/3) \|f\|_{U^4}^4$ on odd-order groups, so $0.5587725 \leq c_4^{\text{odd}} \leq 2/3$ among sources 01–39 (a later placed source records a better upper bound); the prime-torsion onset $h(p) = D_{\text{div}}(p) + 1$, $h(7) = 7$, $h(11) = h(13) = h(17) = 8$, $h(p) \sim 2 \log p / \log \log p$; exact reduction of the characteristic-five cube profile to $\mathbb{Z}/5\mathbb{Z}$, so the u^8 coefficient $20/9$ is universal, with the u^{10} and u^{12} coefficients; an exact six-term norm-extension defect. *Part IV*: the sharp symmetry threshold $1 - (1 - 1/p)^k$ for multilinear selectors over $\mathbb{F}_p^n$ and a lossless polarization identity at Proposition 17.2. *Part V* gains 73 further questions (431 in all). No statement of the new sources is formalized.

[*Eighth write, 7 October 2026: sources 40–46.*] Seven further manuscripts, of 6 and (by their title pages) 7 October 2026, are added to the same Parts, so the report now merges forty-six. *Part I*: a written proof of the whole of Corollary 5.8 in every degree and over every cyclic group with no scale premise: gain $\alpha/3$ at length $\max\{1, (N/M)^{1/D_k}/8\}$, $D_1 = 4$, $D_k = (k!)^2 2^{(k+1)^2}/4 + 1$, through an explicit all-degree recurrence and a lossless recursive localization (during this write the repository closed the formal statement, with the catalogue's weaker constants, along the same architecture and its own finite localization); and an arithmetic

proper-progression increment $\delta + (1 - \delta)/L \cdot \max\{q - L, 0\}/(q - 1)$ on every cyclic group, so the variance gain $(1 - \delta)/(2L)$ needs only $N \geq 2(L - 1)^2 + 1$ instead of $N \geq L^3$, with sharp coefficient $1 - 1/C$ under $N \geq CL^2$. *Part II*: one Fejér kernel in place of the repeated Riesz products of Lemma 15.5, retention $\alpha^s \eta^{s-1}/(3^{s-1}(s+1)^s)$, with a kernel-class barrier; the second Whitney coefficient of the cross-section arrangement of Lemma 15.4 for every d , $b_2(8) = 3,280,234,631,541$, and exact core-avoidance products in every characteristic; the three largest additive energies of torsion-free m -sets, and source 14's coset envelope extended from $\varepsilon \leq 1/100$ to $\varepsilon < 1/9$ with all equality cases. *Part III*: $\|f\|_{U^4}^{16} \geq (4795/832) \|f\|_{U^2}^{16}$ for real centered f on odd-order groups, so $0.5587725 \leq c_4^{\text{odd}} \leq 0.6454070$ among sources 01–46 (a later placed source records a better upper bound). *Part IV*: for symmetric trilinear forms over $\mathbb{F}_2^n$ the exact twisted cubic norm $2^{-r/2}$, r the rank of $T(a, a, b) + T(a, b, b)$, in two independent proofs, and Tidor's codimension $8 \log_2(1/\delta)$ improved to $2 \log_2(1/\delta)$; a sharp Boolean certification threshold $1 - 2^{1-d}$; division-free orthogonality on every finite abelian group; the quartic value $11/16$ on $\mathbb{F}_2^2$. *Part V* gains 62 further questions (493 in all). No statement of the new sources is formalized as stated; the formal proof of Corollary 5.8 is described in Section V.1.

[*Ninth write, 7 October 2026: sources 47–53.*] Seven further manuscripts, of 6 and (by their title pages) 7 October 2026, are added to the same Parts, so the report now merges fifty-three. *Part II*: a collision-aware derandomization of the restriction step of Lemma 9.3 (exact group-algebra inventories of energies with repeated labels, deterministic rounding in $O(s^4 n D)$ operations, prescribed cardinality, fixed linear systems, and a fully rational constructive form of source 17's Fejér restriction with its retention); the two-cube degeneracy $\delta_{k,1}(p)$ with second coefficient $-k(3k^3 + 14k^2 + 15k - 14)/(6p^2)$ in every dimension, exact formulas for $k \leq 3$ at every odd prime ($k = 3$ by a finite certificate), an exact count in characteristic three in every dimension, and eventual Laurent polynomiality; a new Section 13 chapter: an explicit threshold $N_0(a_0) = \lceil (8/c(a_0))^{1/\gamma(a_0)} \rceil$ for common-step square extraction uniformly over actual densities $a \geq a_0$, with side at least $N^{3\gamma(a)}$; full-budget witnesses whose retained Stage 13.9 data contain no common-step square of side two, so the older grid interface is not implied, with a repair by reselection; and composite-modulus counterexamples to the two preserved all-moduli encodings of Stage 13.7; at Lemma 16.10, sharp branch-cover laws and a written counterexample to the catalogue's two-premise `lemma_16_10` (formally refuted in the repository since), and a conditional acyclic replacement for Corollary 16.11 from seven stated interfaces. *Part III*: $\|f\|_{U^4}^{16} \geq (29/4) \|f\|_{U^2}^{16}$ and $\|f\|_{U^5}^4 \geq (50/29) \|f\|_{U^2}^4$ for real centered f on odd-order groups, so $0.5587725 \leq c_4^{\text{odd}} \leq 0.6094183$ among sources 01–53 and $c_5^{\text{odd}} \leq 29/50$. *Part IV*: the canonical Boolean tensor C_d has maximum selected energy exactly $1 - (d + 1)/2^d$ in every degree $d \geq 3$ and dimension $n \geq 2$, which answers source 43's quartic question, and the sharp symmetric nonintegrability thresholds are $1/2, 11/16, 13/16, 57/64$ in degrees three to six. *Part V* gains 57 further questions (550 in all). No statement of the new sources is formalized as stated; the repository's formal refutation of `lemma_16_10`, its endpoint characterization and its explicit Section 13 threshold are described in Section V.1.

[*Tenth write, 7 October 2026: sources 54–59.*] Six further manuscripts, of 6 and (by their title pages) 7 October 2026, are added to Parts II and V, so the report now merges fifty-nine. *Part II*: at the Proposition 6.1 endpoint, sharp second-order stability for bounded functions at the polynomial-phase endpoint of every U^d , $D_d(f) \leq (2/q)\varepsilon + b_d\varepsilon^2 + O_d(\varepsilon^3)$ with $q = 2^d$ and $b_d = (3^d - 1)/(2q^2)$ on odd, $(q^2 - 1)/(3q^2)$ on even groups, both sharp (at $d = 3$ the coefficients $13/64$ and $21/64$ of sources 38 and 31), and explicit near-maximal frequency extraction at every order with a cubic inverse bound; source 11's local affine cubic $4t - 10t^2 + 6t^3$ on $t < 4/9$ ($t < 1/2$ without 3-torsion), the exact non-affine energy maximum from $N \geq 16$, and a relative affine repair on Fourier-uniform domains with error $13\varepsilon/50$; exact character-separation ranks $\lfloor \log_2(q + 1) \rfloor$ and $\lfloor \log_3(2q + 1) \rfloor$ and Freiman restrictions polynomial in the doubling constant for every abelian target; at Proposition 8.1 one quadratic phase on every translate with coefficient at least $27/4$, the exact offset ceilings $1, 1/2, 1/3, 1/4$ and sharp stability exponents; Lemma 9.3 on every cyclic modulus with retention exponent 54, and the optimal leading exponent $m \log m / (2 \log a)$ of positive Fourier filters; a written proof of Theorem 13.12 with the printed length and density $(\alpha/2)^{2^{41}}$ through source 40's Fejér selection; bounded Schur deficits as integer partitions and

the complete third torsion-free energy level. *Part V* gains 63 further questions (613 in all). No statement of the new sources is formalized as stated; the repository has since proved **theorem_13_12** formally along source 57’s route, with density exponent 2^{42} (Section V.1).

[*Eleventh write, 7 October 2026: sources 60–69.*] Ten further manuscripts, of 6 and 7 October 2026, are added to all five Parts, so the report now merges sixty-nine. *Part I*: an exact formula for the missing-arc mass under unequal Fourier budgets near source 37’s equal-budget branch. *Part II*: source 11’s affine cubic $4t - 10t^2 + 6t^3$ on the whole range $t \leq 2/3$ for maps into elementary abelian 3-groups, with the sharp unique-recovery thresholds $1/4$ and $4/9$ and the non-affine energy maximum from $N \geq 13$, optimal; the energy gap $m r_A(d)(m - r_A(d)) \leq m^3 - E(A)$ with coefficient one, the coset envelope through $\varepsilon \leq 1/4$ and uniqueness below $2/9$, both optimal; bounded Schur defects as partitions for $n \geq 3r + 1$ and the fourth torsion-free energy level; Lemma 9.3 on every cyclic group with exponent 53 for arbitrary abelian targets, with retention $\alpha^{31}\eta^{30}$ for cyclic targets, and threshold-free exact Freiman restriction on arbitrary finite abelian groups; source 17’s retention at a size threshold one power of L lower on prime moduli; the complete arrangement polynomial of every order with all equality cases and optimal inverse coefficients, and sharp two-direction repair; the exact interpolation count $E_t(L_1, \dots, L_r)$ at Lemma 16.10. *Part III*: $\|f\|_{U^2}^4 \leq (1/3 - 4/(3N^2) + 1/N^3) \|f\|_{\infty}^4$ on odd-order groups, sharp, hence $\lim_d c_d^{\text{odd}} = 1/3$ with an explicit rate; binomial separation and endpoint bounds $\rho_4(\alpha) \ll \alpha^5$, $\rho_6(\alpha) \ll \alpha^{7.847}$ at Gowers’s Section 4 conjectures; independent norm layers classified with no degree cap. *Part IV*: the largest energy of a nonintegrable symmetric d -linear tensor on $\mathbb{F}_2^n$ is exactly $1 - (d + 1)/2^d$ in every degree, proved twice; the local phase-removal coefficient at Proposition 17.7 is asymptotic to its ceiling $2^{k+1}/(k + 2)!$. *Part V* gains 104 further questions (717 in all). No statement of the new sources is formalized as stated; the repository has since formalized the affine case of source 67’s interpolation count in its own form, and has repaired the catalogue’s Lemma 16.10 to a cover whose graph count has the same compressed form, with a companion that is counted but not yet kernel-checked (Section V.1).

Contents

F.1	Guide to this report	24
F.2	What is proved, and where	45
F.3	Notation across the five sources	62
F.4	Provenance and merge decisions	66
F.5	Relation to the formal project	82
F.6	Source 01: The two interfaces and the main conclusions	86
F.7	Source 02: Scope, sources, and the precise improvements	88
F.8	Source 03: Scope, starting point, and principal results	90
F.9	Source 04: Purpose, results, and scope	92
F.10	Source 05: Scope, baseline, and main contributions	94
F.11	Source 06: Scope, conclusions, and mathematical provenance	96
F.12	Source 07: The problem and the concrete improvements	99
F.13	Source 08: The target lemma and the contribution	100
F.14	Source 09: Scope, source theorem, and principal results	102
F.15	Source 10: Scope, results, and the baseline	104
F.16	Source 11: The problem and the principal results	107
F.17	Source 12: Purpose, main results, and relation to the sources	110
F.18	Source 13: Scope, results, and the source baseline	113
F.19	Source 14: Scope, contributions, and relation to the sources	115
F.20	Source 15: Results, scope, and the source discrepancy	118
F.21	Source 16: Scope, source comparison, and principal results	121
F.22	Source 17: The target and the main conclusions	124
F.23	Source 18: Results, scope, and their relation to the source	127
F.24	Source 19: Scope and principal results	131

F.25	Source 20: The target and the main conclusions	134
F.26	Source 21: The question and the contribution	136
F.27	Source 22: Purpose, results, and scope	139
F.28	Source 23: Results, provenance, and their place in the proof	142
F.29	Source 24: Scope, headline results, and comparison with the source	145
F.30	Source 25: Scope, main theorem, and relation to the source	148
F.31	Source 26: The question and the scope	150
F.32	Source 27: The target and the main conclusions	151
F.33	Source 28: Results and the scope of the improvement	155
F.34	Source 29: Purpose, contributions, and their precise scope	158
F.35	Source 30: Target, contributions, and comparison standard	161
F.36	Source 31: The problem, the contributions, and their scope	163
F.37	Source 32: The theorem and its scope	166
F.38	Source 33: The problem and the contribution	168
F.39	Source 34: Scope, principal results, and the audited baseline	171
F.40	Source 35: The research target and the principal conclusions	174
F.41	Source 36: The problem and the main result	178
F.42	Source 37: Research target and relation to the repository	180
F.43	Source 38: Research question, main results, and provenance	183
F.44	Source 39: The local problem and the scope of the improvement	187
F.45	Source 40: The local problem and the contribution	190
F.46	Source 41: The theorem and its scope	192
F.47	Source 42: The arithmetic increment and its meaning	194
F.48	Source 43: The research problem and the principal results	196
F.49	Source 44: The research target and what is new	200
F.50	Source 45: Purpose, results, and provenance	202
F.51	Source 46: Research target, contribution, and boundaries	206
F.52	Source 47: The problem and the main results	208
F.53	Source 48: The research target and the resulting strengthening	211
F.54	Source 49: Scope, baseline, and principal contributions	214
F.55	Source 50: The obstruction and its significance	219
F.56	Source 51: The conclusion and its quantifiers	221
F.57	Source 52: Main conclusions and source scope	224
F.58	Source 53: The explicit theorem	226
F.59	Source 54: The result and its place in the proof	229
F.60	Source 55: The endpoint problem and its place in Gowers’s proof	232
F.61	Source 56: The question and the main result	234
F.62	Source 57: Principal results and their scope	237
F.63	Source 58: Main results and their place in the existing work	241
F.64	Source 59: Scope, main results, and relation to the source	245
F.65	Source 60: Research context and principal conclusions	249
F.66	Source 61: Research objective, source comparison, and main results	254
F.67	Source 62: The target and the contribution	257
F.68	Source 63: Introduction and principal results	258
F.69	Source 64: The problem, the results, and their position in the literature	262
F.70	Source 65: The extremal problem and the main results	264
F.71	Source 66: Scope, previous results, and principal contributions	268
F.72	Source 67: Research position and the exact improvements	270
F.73	Source 68: Introduction and contribution map	275
F.74	Source 69: The problem, the results, and the comparison baseline	279

I	Density transfer and phase-flat partitions	283
I.1	The finite weighted phase problem [01]	287

I.2	A sharp density-tail theorem [01]	290
I.3	The optimal forced maximum and its tradeoff [01]	292
I.4	Partition discrepancy and finite convex optimization [01]	293
I.5	Integer-exact simultaneous progression partitions [01]	296
I.6	A strengthened Fourier density-increment lemma [01]	297
I.7	A sharp scale barrier for the phase-flat method [01]	300
I.8	Sharp phase freezing for real functions [02]	301
I.9	A sharp density-mass tradeoff [02]	303
I.10	Combining correlation, phase control, and density extraction [02]	305
I.11	A uniform refinement theorem and an optimized phase budget [02]	306
I.12	A rational equality example for the moment bound [02]	309
I.13	Conventions and analytic background [03]	309
I.14	An exact selection theorem for density increments [03]	310
I.15	A density-sensitive Fourier increment on proper progressions [03]	312
I.16	Exact finite partitions and the sharp rounding constant [05]	314
I.17	Density-sensitive localization and popular increments [05]	317
I.18	Simultaneous phases and an optimal exponent [05]	320
I.19	Spectral energy, residual variance, and unequal tolerances [05]	323
I.20	Average-length partitions and exceptional-cell budgets [05]	326
I.21	Sharp density extraction from a partition [06]	327
I.22	Conventions and normalization [10]	331
I.23	Sharp density profiles on a partition [10]	332
I.24	A constructive Fourier-to-progression tradeoff [10]	334
I.25	Conventions and normalization [13]	337
I.26	Simultaneous linear phases: finite partitions and optimal powers [13]	339
I.27	Sharp selection from partition imbalance [13]	343
I.28	Real phase removal and a joint extremal theorem [13]	345
I.29	A density-sensitive Fourier increment [13]	348
I.30	Direct partitions for affine frequencies and their own steps [16]	351
I.31	An obstruction and a precise remaining partition problem [16]	355
I.32	Simultaneous polynomial partitions by descent in degree [18]	356
I.33	A volume obstruction to polynomial partition exponents [18]	363
I.34	Simultaneous polynomial phase partitions [19]	366
I.35	The partition problem and its arithmetic parameter [23]	374
I.36	An explicit theorem with matching upper and lower scales [23]	375
I.37	Sharp asymptotics and the permitted range of cell sizes [23]	378
I.38	Consequences for Lemma 5.9 and exact finite witnesses [23]	380
I.39	Partial coverage and the exact cost of exceptions [23]	382
I.40	Corollary 5.8 without a scale threshold: degrees one and two [write]	384
I.41	Definitions and the two main bounds [26]	386
I.42	Two elementary mechanisms [26]	388
I.43	Proofs of the main theorems [26]	391
I.44	Common step integer box rectification [26]	393
I.45	An explicit quadratic recurrence estimate [32]	395
I.46	Lossless phase partitions [32]	398
I.47	A second moment over every cyclic modulus [32]	400
I.48	Assembly and every scale [32]	401
I.49	Source 37's chapter: missing intervals at Lemmas 5.2 and 5.5 [write]	403
I.50	The exact gap-moment theorem [37]	403
I.51	A positive-coefficient polynomial certificate [37]	405
I.52	A supported trigonometric moment criterion [37]	407
I.53	Attainment, support, and equality [37]	408
I.54	Endpoints and the first spectral transition [37]	410
I.55	A linear-frequency replacement for Lemma 5.2 [37]	412

I.56	Stability and near extremizers [37]	414
I.57	A stronger monomial-recurrence deduction [37]	416
I.58	The real interval factorization used above [37]	420
I.59	Low-dimensional examples [37]	421
I.60	Source 35's Part 3: polynomial divisibility and prescribed lengths [write]	422
I.61	Polynomial divisibility and prescribed progression lengths [35]	422
I.62	Sources 41 and 42: Corollary 5.8 in every degree, and the quadratic variance scale [write] . .	430
I.63	An explicit quadratic recurrence estimate [41]	431
I.64	Explicit recurrence in every degree [41]	434
I.65	Lossless phase partitions [41]	438
I.66	A second moment over every cyclic modulus [41]	441
I.67	Assembly at every scale [41]	443
I.68	Exact budgets and sharper alternatives [41]	444
I.69	The covariance square and the range inequality [42]	449
I.70	Exact support optimization and integer scales [42]	451
I.71	Positive definite weights and the classical packing bound [42]	453
I.72	Exact extremizers and the optimal uniform constant [42]	455
I.73	A density sensitive gain at the quadratic endpoint [42]	456
I.74	A sparse affine cover consequence [42]	458
I.75	Source 67's Section 2: unequal Fourier budgets at Lemmas 5.2 and 5.5 [write]	461
I.76	Exact Fourier-gap profiles with unequal budgets [67]	462

II The inverse step: energies, moments, restriction and dependent random choice 470

II.1	Normalization and finite Fourier identities [02]	479
II.2	A sharp weighted extension of Proposition 6.1 [02]	480
II.3	The L^2 denominator and equality examples [02]	482
II.4	Keeping the full derivative spectrum and tracking rigidity [02]	484
II.5	Sharper common-neighborhood selection [03]	486
II.6	An elementary quantitative replacement for Proposition 7.3 [03]	489
II.7	Notation and normalization [04]	492
II.8	Exact second moments and density-sensitive lifting [04]	493
II.9	Polynomial restriction to approximate homomorphisms [04]	496
II.10	Arrangement moments retain the cross-section profile [04]	502
II.11	Source 07's chapter: transport rigidity at Lemma 10.9 [write]	505
II.12	Notation and assumptions at a glance [07]	506
II.13	Finite fibered domains and the correct error law [07]	506
II.14	Transport profiles instead of square-root losses [07]	507
II.15	The coupled-path rigidity theorem [07]	510
II.16	The balanced threshold is sharp [07]	512
II.17	Application to Gowers's Section 10 [07]	513
II.18	Choosing the initial law by a rational linear program [07]	517
II.19	Further proved extensions [07]	519
II.20	An explicit weighted replacement for the energy extraction step [10]	521
II.21	A sharper obstruction-fibre form of the Freiman extraction [10]	523
II.22	Source 11's chapter: affine rigidity at the Section 6 endpoint [write]	528
II.23	Graph energy, derivatives, and the exact endpoint [11]	528
II.24	Exact counting for a one-valued error [11]	530
II.25	Replacing many error values by one: odd targets [11]	531
II.26	Targets with two-torsion, including mixed targets [11]	532
II.27	Quantitative information about near-extremizers [11]	533
II.28	Removing the assumption of proximity to an affine map [11]	534
II.29	Exact inversion and sharp recovery constants [11]	535

II.30	The largest non-affine energy in arbitrary large finite groups [11]	536
II.31	All finite vector-space sizes: a second proof [11]	537
II.32	Partial domains and the interface with Proposition 6.1 [11]	539
II.33	Source 12's chapter: Freiman extension and Bohr radii in Section 10 [write]	541
II.34	Translation loss and intersections along a tree [12]	542
II.35	A sharp theorem for induced Freiman homomorphisms [12]	543
II.36	Correction from concentrated directional increments [12]	545
II.37	A finite selection argument for Bohr radii [12]	547
II.38	Fourier smoothing with the full fourth moment [12]	549
II.39	Synchronizing a local model across fibres [12]	551
II.40	A strengthened version of the Section 10 local-model theorem [12]	552
II.41	Consequences and the size of the improvement [12]	556
II.42	The fixed-domain repair in the Section 10 extraction input [12]	558
II.43	Sharp stability of the cross-section collision bound [14]	559
II.44	Equality and stability in arrangement moment interpolation [14]	566
II.45	Source 15's chapter: collision-sensitive amplification at Lemmas 14.6–14.8 [write]	570
II.46	Axis cubes and collision measures [15]	571
II.47	Exact moment amplification [15]	572
II.48	The cubic fibre bound and sharp density improvements [15]	573
II.49	Coupling the product property to the collision count [15]	574
II.50	Quantitative comparison with the supplied statements [15]	578
II.51	Equality in the collision cap [15]	579
II.52	Dimension-free Cartesian stability [15]	581
II.53	Stability of the amplification inequality [15]	583
II.54	A complete normal form for joint equality [15]	585
II.55	A geometry–amplification dichotomy [15]	587
II.56	Vertex weights and other extensions [15]	588
II.57	Source 17's chapter: polynomial-loss restriction at Section 9 [write]	589
II.58	Notation and elementary counting facts [17]	590
II.59	Moment lifting without a density loss [17]	591
II.60	A single Fejér filter [17]	592
II.61	The restriction theorem with all exceptional tuples included [17]	594
II.62	A limitation of the scalar filter-certificate method [17]	597
II.63	Fixed characteristic: arbitrary linear systems and rank profiles [17]	598
II.64	Boundary cases, normalization traps, and constructive content [17]	601
II.65	Sources 18 and 19: degenerate arrangements (Section 15) and common sampling (Section 16) [write]	602
II.66	A sharp first-order count of degenerate arrangements [18]	603
II.67	The finite common-sampling problem [19]	612
II.68	An exact minimax theorem and a deterministic algorithm [19]	615
II.69	Sharp asymptotics of the sampling loss [19]	617
II.70	Recovering polynomial pieces from a common set of slices [19]	620
II.71	Quantitative interpretation and worked comparisons [19]	626
II.72	Source 24's chapter: Section 7 again [write]	628
II.73	Notation and the exact moment–defect principle [24]	629
II.74	A three-sample strengthening of Lemma 7.4 [24]	631
II.75	The sharp scalar relaxation and a two-sample barrier [24]	632
II.76	From a profile certificate to many four-step walks [24]	634
II.77	An explicit energy-to-small-difference-set theorem [24]	635
II.78	The sunset-growth input, with a complete proof [24]	637
II.79	Vertical-fibre compression and Freiman restrictions [24]	638
II.80	Compiling energy bounds into coherent restrictions [24]	641
II.81	Sources 22, 25 and 27 on Lemma 15.4, and source 23 on Lemma 16.10 [write]	643
II.82	Arrangements, parity, and the degree of a certificate [22]	645

II.83	An exact polynomial count for the first nonparity degree [22]	647
II.84	The finite exceptional-set bound and its quantitative use [22]	650
II.85	The exceptional two-cube case: a signed-coordinate arrangement [22]	654
II.86	Arrangements and their moment relations [25]	657
II.87	Minimal Walsh layers give independent linear equations [25]	658
II.88	Non-parity witnesses have codimension at least two [25]	659
II.89	The cross-section hyperplanes and their exact number [25]	661
II.90	The global count and the geometry of the exceptional set [25]	663
II.91	Explicit ambient-size savings [25]	665
II.92	A finite-field random-selection theorem with exact parameters [25]	666
II.93	Labelled arrangements and their moment relations [27]	668
II.94	Boolean coefficient extraction and a rank lemma [27]	670
II.95	Counting coefficient directions exactly [27]	671
II.96	All first-order obstructions [27]	673
II.97	Second-order structure on the regular region [27]	675
II.98	Two pairs of cross sections and the second coefficient [27]	678
II.99	Two extensions of the second-order theorem [27]	679
II.100	An exact count for eight labelled vertices [27]	681
II.101	Consequences for an exceptional-arrangement budget [27]	685
II.102	Boundary cases and safeguards [27]	686
II.103	Common base partitions and exact transport of affine fibres [23]	687
II.104	Sources 31 and 29: the Proposition 6.1 endpoint [write]	694
II.105	Normalized derivatives and their phase space [31]	695
II.106	Recovering an extremal subgroup without an exponent loss [31]	697
II.107	The sharp local quartic envelope [31]	700
II.108	Sharp stability of a selected derivative graph [31]	703
II.109	The correct stability statement for the weights [31]	704
II.110	A joint endpoint rigidity theorem [29]	707
II.111	Soft selectors [29]	712
II.112	Mixed endpoint rigidity and affine derivative selectors [29]	714
II.113	Sharpness of the error order [29]	715
II.114	Source 31's chapter: weights in the Section 7 restriction [write]	716
II.115	An elementary weighted energy extraction [31]	716
II.116	A weighted finite form of Gowers' Freiman restriction [31]	718
II.117	Keeping the original derivative correlation [31]	720
II.118	A square-root weighted mass theorem from Reiher-Schoen [31]	721
II.119	Source 29 at Lemma 9.3: restriction over a fixed field [write]	723
II.120	Restriction without discarding affine degeneracies [29]	723
II.121	Sources 28 and 30 at Section 10: logarithmic selection and fixed radii [write]	729
II.122	Conventions and finite domains [28]	730
II.123	Weighted local errors and a useful anchor [28]	730
II.124	Logarithmic component selection [28]	732
II.125	From an averaged anchor error to exact local differences [28]	735
II.126	Fourier smoothing and the proof of the main theorem [28]	740
II.127	Consequences, comparisons, and structural limits [28]	742
II.128	Conventions and finite translation calculus [30]	744
II.129	The escape-tube inequality [30]	745
II.130	Packing the enlarged Bohr set [30]	747
II.131	Dense-domain extension of partial homomorphisms [30]	748
II.132	Larger Bohr neighborhoods in Corollary 10.11 [30]	750
II.133	A density-sensitive compatibility step for Lemma 10.12 [30]	751
II.134	Further consequences without radius regularization [30]	753
II.135	Exponential-rank obstructions, including cyclic groups [30]	754
II.136	Source 33's chapter: high-agreement repair of respected arrangements [write]	757

II.137	Sampling, independence, and Fourier normalization [33]	757
II.138	A noisy homomorphism lemma with auxiliary labels [33]	760
II.139	Synchronizing derivatives into one biadditive model [33]	762
II.140	Sharp local defect inequalities [33]	763
II.141	Extremal families and optimal coefficients [33]	766
II.142	Erasures and almost-full products [33]	769
II.143	Exact cyclic decoding and reproducible checks [33]	770
II.144	Source 35's Part 1: degree-weighted selection at Lemma 15.5 [write]	772
II.145	The arrangement model and the quantitative theorem [35]	773
II.146	Degree balances and the probability of a certificate [35]	775
II.147	The random restriction and its collision correction [35]	778
II.148	The exponent, a fair comparison, and a downstream substitution [35]	781
II.149	Source 38's Part II: sharp odd-order stability at the Proposition 6.1 endpoint [write]	784
II.150	The sharp odd-order local stability coefficients [38]	785
II.151	Source 39's chapter: shared anchors at Lemma 16.10 [write]	795
II.152	A finite model of shared-anchor certification [39]	795
II.153	A cutoff-free shared sampling theorem [39]	796
II.154	The exact finite minimax problem [39]	798
II.155	A closed finite profile formula and the affine quadratic cap [39]	799
II.156	The Poisson profile and its exact optimization [39]	803
II.157	Uniform asymptotics and a sharp sample-complexity law [39]	804
II.158	Deterministic anchor selection by conditional expectation [39]	808
II.159	Polynomial reconstruction, general interpolation spaces, and errors [39]	809
II.160	Labelled slice covers and candidate counts [39]	812
II.161	Hereditary refinement without loss of slice labels [39]	813
II.162	Source 34's Appendix A: the sharp nonuniform transport threshold [write]	815
II.163	The sharp nonuniform directional rigidity threshold [34]	815
II.164	Source 40's chapter: one-stage Fejér selection at Lemma 15.5 [write]	817
II.165	Arrangements, respect, and the counting normalization [40]	817
II.166	A single monomial isolates the universal relation [40]	819
II.167	Bounded-coefficient exceptional arrangements [40]	820
II.168	Fejér kernels and exact survival probabilities [40]	821
II.169	The one-stage selection theorem [40]	823
II.170	Comparison with Gowers' Lemma 15.5 [40]	825
II.171	Vector targets and exact torsion dependence [40]	827
II.172	A sharp exponent obstruction for positive Fourier kernels [40]	829
II.173	Source 44's chapter: the grid core of Lemma 15.4 [write]	830
II.174	Arrangements, coefficient directions, and the regular region [44]	831
II.175	A width theorem with no dimension-dependent characteristic cutoff [44]	832
II.176	Exact core-avoidance laws [44]	833
II.177	From exact core avoidance to the full degeneracy count [44]	835
II.178	The second Whitney coefficient and a uniform characteristic cutoff [44]	838
II.179	A fixed nine-point grid computes every second coefficient [44]	839
II.180	Eight lattice energies and the explicit eight-arrangement coefficient [44]	841
II.181	The exponentially smaller pair-overlap deficit [44]	844
II.182	Contrasting growing-dimension laws [44]	846
II.183	Source 45's chapter: energy levels and the coset range [write]	847
II.184	The first three additive-energy levels [45]	848
II.185	A larger coset-stability range and the geometry of equality [45]	855
II.186	Source 48's chapter: collision-aware derandomization of the restriction step [write]	867
II.187	Sampling counts and the repeated-entry obstruction [48]	867
II.188	An exact group-algebra inventory [48]	868
II.189	Local replacement, Stirling coefficients, and bit cost [48]	870
II.190	Deterministic signed-energy rounding [48]	872

II.191	Exact cardinality without a cardinality-sized state [48]	873
II.192	General linear systems and column-type compression [48]	876
II.193	Robust certificates and entirely rational filters [48]	877
II.194	A deterministic polynomial-loss restriction theorem [48]	879
II.195	Sources 53, 51 and 52: Section 13 square geometry and Stage 13.7 [write]	883
II.196	The context and the quantifiers [53]	884
II.197	Scalar parameters and imported analytic inputs [53]	885
II.198	Seven sufficient quantitative budgets [53]	887
II.199	The seventh threshold dominates exactly [53]	889
II.200	Why the all-density witness is short [53]	892
II.201	Short parents and direct square geometry [53]	892
II.202	The sharper coefficient and final assembly [53]	895
II.203	The exact retained interface [51]	896
II.204	The geometric obstruction in isolation [51]	898
II.205	Uniform constants and a simultaneous prime choice [51]	898
II.206	A genuine strip context and the initial data [51]	901
II.207	Verification of all six budgets [51]	902
II.208	The complete chosen output chain [51]	903
II.209	Failure of the grid and of the new fit hypothesis [51]	905
II.210	A conditional covering repair and its exact scope [51]	905
II.211	A direct repair for suitably selected witnesses [51]	910
II.212	The precise retained interface [52]	914
II.213	The common dyadic strip [52]	917
II.214	The exact step two counterexample [52]	918
II.215	Certifying both genuine initial stages [52]	919
II.216	The obstruction persists with unit input step [52]	922
II.217	Genuine initial stages with all six numerical budgets [52]	923
II.218	A prime power family of unit step obstructions [52]	927
II.219	A small model and the algebraic boundary [52]	928
II.220	Source 49's chapter: exact two-cube degeneracy [write]	929
II.221	Two-cube degeneracy as a translation problem [49]	930
II.222	Rank-two structure and the second degeneracy coefficient [49]	933
II.223	Exactness in dimension three: a finite proof for all odd primes [49]	936
II.224	Eventual polynomiality in every fixed dimension [49]	939
II.225	Finite-field extensions and an exact subfield-alphabet theorem [49]	941
II.226	Sources 50 and 49: Lemma 16.10 is false as catalogued, and an acyclic replacement [write]	943
II.227	Progressions, polynomial lists, and branch masses [50]	944
II.228	Exact one-dimensional cover laws [50]	945
II.229	Product boxes, adaptive partitions, and global rigidity [50]	947
II.230	Balanced words at logarithmic scales [50]	949
II.231	Realizing arbitrary local cover profiles [50]	951
II.232	The pinned Section 16 interface [50]	953
II.233	A counterexample to the two-input form of Lemma 16.10 [50]	955
II.234	A repair interface that retains complexity [50]	959
II.235	An acyclic replacement for the Section 16 cover theorem [49]	962
II.236	Sources 55 and 57: the polynomial-phase endpoint in every degree [write]	975
II.237	Main results [55]	976
II.238	Demodulation and the bounded-function geometry [55]	978
II.239	Four vertices of a discrete cube [55]	979
II.240	A cubic remainder for bounded complex perturbations [55]	981
II.241	A continuous two-torsion profile [55]	983
II.242	The explicit local stability theorem [55]	984
II.243	Uniform entry and matching extremizing families [55]	986
II.244	Proof of second-order rigidity [55]	988

II.245	An exact calibration at degree two [55]	990
II.246	Consequences for phase testing and local structure [55]	990
II.247	Near-maximal uniformity and global frequency maps [57]	991
II.248	Sources 59 and 56: affine rigidity on larger ranges and relative repair [write]	999
II.249	Signed graph expansions and stronger affine rigidity [59]	1000
II.250	Averaging the signed kernel: radii four ninths and one half [59]	1003
II.251	Classical correction and global affine recovery [59]	1007
II.252	Examples, comparisons, and precise claim boundaries [59]	1010
II.253	Normalization and elementary Fourier facts [56]	1012
II.254	Autocorrelation, triangle variance, and pointwise links [56]	1013
II.255	Weighted modal differences without discarding fibers [56]	1015
II.256	Self-correction and coarse affine reconstruction [56]	1016
II.257	The one-error profile and linear amplification [56]	1018
II.258	Proof of the main theorem and optimal first-order coefficient [56]	1020
II.259	Exact extension, spectral certificates, and testing [56]	1022
II.260	Explicit affine repair on quadratic level sets [56]	1023
II.261	Source 59's chapter: exact character separation at Section 7 [write]	1026
II.262	Exact character distributions and order-sensitive radii [59]	1026
II.263	Adaptive character separation and matching obstructions [59]	1028
II.264	Weighted Freiman restrictions over arbitrary abelian targets [59]	1030
II.265	Small doubling and comparison with Lemma 7.5 [59]	1032
II.266	One-character restrictions and lonely runners [59]	1035
II.267	Source 54's chapter: one quadratic phase at Proposition 8.1 [write]	1038
II.268	An exact mixed and weighted Fourier factorization [54]	1039
II.269	Quadratic untwisting and the strengthened proposition [54]	1041
II.270	Every finite abelian group, including even order [54]	1042
II.271	Exact global offset ceilings and equality cases [54]	1043
II.272	Dimension-free stability and the fourth-root obstruction [54]	1045
II.273	Several intercepts and the classical clique parameter [54]	1048
II.274	Localization, robustness, and many useful translates [54]	1049
II.275	An exact deficit and a constructive selector [54]	1051
II.276	Source 58's chapter: Lemma 9.3 on every cyclic group [write]	1052
II.277	Counts, normalizations, and finite Fourier analysis [58]	1053
II.278	Admissible filters and their torsion spectrum [58]	1054
II.279	A finite random selector and the good-tuple bound [58]	1055
II.280	Nonprincipal relations and the domain torsion cost [58]	1056
II.281	Repeated bad tuples and their decaying error [58]	1058
II.282	The expectation theorem and a simultaneous selection certificate [58]	1059
II.283	Powered cosine filters and the certified exponent 54 [58]	1061
II.284	Further generalizations of the restriction estimate [58]	1063
II.285	Optimal leading growth of the Fourier-filter exponent [58]	1065
II.286	Moment lifting and the Gowers-ProveIt interfaces [58]	1068
II.287	Source 57's chapter: Theorem 13.12 in writing, and formally [write]	1070
II.288	Conventions and the exact imported results [57]	1071
II.289	Fejér purification for the Section 13 arrangements [57]	1073
II.290	Improved quantitative propagation and the printed exponent [57]	1075
II.291	Source 57's chapter: bounded Schur deficits and the third energy level [write]	1079
II.292	Bounded Schur defects and integer partitions [57]	1079
II.293	Source 67's chapter: affine rigidity in characteristic three [write]	1086
II.294	Cubic rigidity throughout the nontrivial range in characteristic three [67]	1087
II.295	Sharp thresholds for global affine recovery [67]	1093
II.296	Sources 68 and 66: the energy gap, the coset range and the fourth energy level [write]	1098
II.297	A sharp pointwise energy inequality [68]	1099
II.298	Sharp coset rigidity and two different thresholds [68]	1103

II.299	Sharp puncture repair and a single near-saturated translation [68]	1107
II.300	Exact ordered defects and endpoint reflection [66]	1116
II.301	Two exact reflections force an arithmetic prefix [66]	1118
II.302	The partition normal-form theorem [66]	1118
II.303	Boundary-hole normal forms for endpoint extensions [66]	1120
II.304	An exact energy law for arbitrary boundary holes [66]	1122
II.305	A sharp all-defect endpoint envelope [66]	1123
II.306	The known first and second levels, with complete proofs [66]	1125
II.307	Complete classification of the third energy level [66]	1126
II.308	The fourth energy level and the eight-template theorem [66]	1127
II.309	Inverse consequences and the boundary of the Ramsey application [66]	1129
II.310	Sources 68, 62, 61 and 60: Lemma 9.3 on every modulus [write]	1131
II.311	A restriction estimate with decaying exceptional terms [68]	1132
II.312	A degree-twelve filter and the exponent fifty-three [68]	1137
II.313	Uniform principal costs as arity and torsion vary [68]	1140
II.314	Notation and the main restriction theorem [62]	1147
II.315	The unit-dilation detector [62]	1149
II.316	Uniform primitive-response bounds [62]	1151
II.317	Weighted short relations over composite cyclic groups [62]	1152
II.318	The restriction proof, including repeated vertices [62]	1155
II.319	An explicit order-eight strengthening [62]	1157
II.320	Torsion obstructions and the scope of optimality [62]	1158
II.321	Extensions and constructive content [62]	1160
II.322	Notation and the exact relation interface [61]	1161
II.323	Separating every nonzero defect with few characters [61]	1162
II.324	Small differences and a vertical-defect budget [61]	1164
II.325	Exact Freiman partitions and retention statements [61]	1165
II.326	A self-contained energy-to-growth input [61]	1167
II.327	Polynomial exact restriction of high relation energy [61]	1169
II.328	Consequences for Gowers’s restriction lemmas [61]	1171
II.329	Torsion obstructions and the optimal logarithmic scale [61]	1173
II.330	A translated Fejér restriction theorem [60]	1176
II.331	Source 60’s chapter: the arrangement polynomial of every order [write]	1182
II.332	Respected arrangements and their normalization [60]	1182
II.333	A sharp spectral inequality with independent labels [60]	1184
II.334	The complete polynomial and its equality cases [60]	1186
II.335	A repair modulus sharp to every asymptotic order [60]	1189
II.336	Odd targets and a sharper common-prime coefficient interval [60]	1191
II.337	Two-direction testing and sharp robust intersection [60]	1195
II.338	Source 67’s Section 5: the interpolation count at Lemma 16.10 [write]	1200
II.339	The exact number of interpolation candidates [67]	1201

III Cube and progression counts 1209

III.1	Progression counting with fewer error terms [03]	1214
III.2	The exact fourth-order term in the cube count [03]	1216
III.3	Sharp centered refinements of the cube-counting lemma [04]	1227
III.4	Density-weighted progression counting [04]	1236
III.5	Quadratic discrepancy bounds for arithmetic progressions [06]	1240
III.6	Normalization and the basic analytic tools [08]	1245
III.7	Three exact cancellations and the fourth-order bound [08]	1247
III.8	The exact quartic term and the role of two-torsion [08]	1248
III.9	The complete three-cube expansion in odd order [08]	1251
III.10	A counterexample to the quartic global lower bound [08]	1254

III.11	Sharpness of the fourth-order exponent [08]	1255
III.12	Lower bounds and Fourier stability [08]	1258
III.13	Different weights at different vertices [08]	1259
III.14	Which centered terms can survive? [08]	1260
III.15	A short exact certificate for the counterexample [08]	1261
III.16	Normalized cube forms and the analytic input [09]	1262
III.17	Balancing removes the first three orders [09]	1264
III.18	Integer value lattices and Fourier constraints [09]	1265
III.19	The complete four-vertex classification [09]	1266
III.20	The complete five-vertex classification [09]	1267
III.21	Exact expansion and quantitative consequences [09]	1269
III.22	Sharpness, signs, and transfer to genuine sets [09]	1272
III.23	A sixth-order order-five obstruction and an all-prime family [09]	1275
III.24	Heterogeneous functions and an all-orders organizing principle [09]	1277
III.25	Relation to progression bounds and current conjecture status [09]	1278
III.26	Exact two-frequency polynomial certificate [09]	1279
III.27	An integer determinant certificate for the order-five example [09]	1280
III.28	A torsion-sensitive generalized von Neumann inequality [10]	1280
III.29	Density-sensitive counting and modular progression criteria [10]	1284
III.30	An exact dimension threshold for polynomial colorings [10]	1287
III.31	Centered cube counts and their spectral decomposition [13]	1292
III.32	Interval discrepancy with the sharp length exponent [13]	1297
III.33	Progression counting by telescoping [13]	1299
III.34	The exact secondary term in the three-cube extremal problem [14]	1301
III.35	Order five changes the next exponent [14]	1308
III.36	Exact supersaturation and rigidity for polynomial colorings [14]	1311
III.37	Quantitative examples and what the refinements measure [14]	1317
III.38	A finite certificate for the centered cube expansion [14]	1319
III.39	Exact fourth- and fifth-order cube corrections in every dimension [16]	1320
III.40	Exact fourth and fifth order cube counts [18]	1327
III.41	Source 20's chapter: quadratic counting errors at Section 3 [write]	1335
III.42	Normalized cubes and subgroup directions [20]	1336
III.43	The two-function directional estimate [20]	1337
III.44	Sharp conversion from directional to global norms [20]	1340
III.45	Second-order telescoping and counting [20]	1341
III.46	Pair correlations as an exact torsion correction [20]	1344
III.47	Supersaturation and the uniformity parameter [20]	1345
III.48	A sharp quadratic obstruction for three-term counts [20]	1346
III.49	Source 21's chapter: the exponent-five coefficient, with source 23's second route [write]	1350
III.50	Fourier identities and the sharp quartic deficit [21]	1350
III.51	Higher centered coefficients with a retained Fourier tail [21]	1352
III.52	The characteristic-five resonance and a finite bound [21]	1354
III.53	Sharp uniform optimization and the phase envelope [21]	1355
III.54	Near equality and the scale of rigidity [21]	1357
III.55	The exact next coefficient on the five-point group [21]	1358
III.56	Indicator sharpness and its limitations [21]	1363
III.57	A quantitative projection consequence of cube Cauchy–Schwarz [21]	1365
III.58	Sharp secondary three-cube asymptotics in characteristic five [23]	1366
III.59	A closer expansion of the five-point example [23]	1372
III.60	Source 22's chapter: degree filtrations of polynomial colourings [write]	1373
III.61	Polynomial zero sets and prescribed equality patterns [22]	1373
III.62	Arithmetic-progression specializations [22]	1377
III.63	Complete quadratic classification and counts [22]	1379
III.64	Why independent additive blocks cannot extend the norm coloring [22]	1380

III.65	A quantitative global obstruction to the high-degree norm construction [22]	1382
III.66	Sources 29 and 28: the odd-order constants c_d^{odd} [write]	1384
III.67	Notation and elementary Fourier identities [29]	1385
III.68	Positive autocorrelation moments and higher norm bounds [29]	1386
III.69	Consequences for the unresolved higher norm constants [29]	1389
III.70	A stronger real norm gap and a higher-order obstruction [28]	1392
III.71	A stronger fourth-dimensional norm comparison [34]	1401
III.72	A stronger gap between real centered uniformity norms [35]	1410
III.73	Source 34's chapter: prime torsion in cube expansions [write]	1419
III.74	Fourier constraints and the Boolean row quotient [34]	1420
III.75	The exact determinant reduction [34]	1421
III.76	Exact first cases and finite nonexistence certificates [34]	1423
III.77	An elementary positive detector for every cyclic order [34]	1425
III.78	Consecutive binary determinants with controlled Smith form [34]	1427
III.79	Asymptotically optimal support for cyclic torsion [34]	1432
III.80	Universal positivity and its exact asymptotic cost [34]	1433
III.81	The exact quartic term and its interface with Gowers's lemma [34]	1436
III.82	Source 38's Part I: exact quotient reduction in exponent five [write]	1437
III.83	Fourier identities and the retained energy defect [38]	1438
III.84	Exact reduction to a five-point quotient [38]	1440
III.85	The five-point profile and its analytic expansion [38]	1446
III.86	Cyclic reduction for every fixed odd exponent [38]	1453
III.87	Source 38's Part III: six-term norm-extension defects [write]	1455
III.88	An exact six-term extension defect and its quadratic-character term [38]	1456
III.89	Source 45's chapter: the real fourth-uniformity gap [write]	1462
III.90	A quantitative real fourth-uniformity gap [45]	1462
III.91	Applications to cube counting and structured holes [45]	1474
III.92	Source 49's chapter: two-harmonic dual tests and $c_4^{\text{odd}} \leq (4/29)^{1/4}$ [write]	1477
III.93	Two-harmonic dual tests and stronger real uniformity gaps [49]	1478
III.94	Sources 63 and 60: the limit of the odd-order constants [write]	1486
III.95	Conventions and standard tools [63]	1486
III.96	A sharp bound for bounded real functions [63]	1487
III.97	The exact high-order odd-group comparison constant [63]	1490
III.98	An asymptotically sharp quartic cube coefficient [63]	1499
III.99	The limiting real Gowers norm constant [60]	1502
III.100	Source 64's chapter: binomial separation at Gowers's Section 4 [write]	1513
III.101	Weighted separation and the six-variable construction [64]	1514
III.102	Why a perfect cyclic six-variable alphabet is impossible [64]	1517
III.103	An exact nine-copy circle transfer [64]	1518
III.104	Exact binomial transfer to the two-torus [64]	1519
III.105	Endpoint exponents from a prime-indexed family [64]	1521
III.106	Deterministic realization in cyclic groups [64]	1523
III.107	An exact dimension barrier for bounded-degree polynomial colors [64]	1526
III.108	Source 67's Section 6: norm layers without a degree cap [write]	1527
III.109	A degree-free obstruction for independent norm layers [67]	1528

IV Floor patterns and the Proposition 17.7 phase extraction 1536

IV.1	Fractional carry patterns and exact range dependence [01]	1539
IV.2	Lower bounds and the optimal entropy scale [01]	1541
IV.3	Affine patterns, translation classes, and general features [01]	1544
IV.4	A polynomial generalization on the Boolean cube [01]	1546
IV.5	Exact finite enumeration and reproducible checks [01]	1548
IV.6	Floor patterns and efficient extraction of a progression profile [04]	1550

IV.7	Exact floor-pattern counts and asymptotically sharp carry entropy [06]	1557
IV.8	A quadratic-exponent loss in localized phase removal [06]	1561
IV.9	Exact computations and reproducible verification [06]	1564
IV.10	Source 16's chapter: Proposition 17.7 without carry patterns [write]	1566
IV.11	Cube forms, normalizations, and finite partitions [16]	1567
IV.12	Explicit polarization and its algebraic limits [16]	1568
IV.13	Finite sliding-window localization without carry patterns [16]	1570
IV.14	A refinement when the original small cells must be kept [16]	1575
IV.15	Source 36's chapter: sharp symmetry certification at Proposition 17.2 [write]	1577
IV.16	Normalization and the exact cube identity [36]	1577
IV.17	The sharp bilinear estimate [36]	1578
IV.18	The derivative rank profile [36]	1580
IV.19	Proof and sharpness of the symmetry threshold [36]	1581
IV.20	Equality structure and stability of the bilinear step [36]	1583
IV.21	Lossless polarization and the Section 17 interface [36]	1585
IV.22	Robust and local variants [36]	1586
IV.23	Multiaffine selectors and extension fields [36]	1588
IV.24	Small characteristic and other exact boundaries [36]	1590
IV.25	Sources 46 and 43: exact binary cubic norms and Boolean phase integration [write]	1591
IV.26	Conventions and the first integration obstruction [46]	1592
IV.27	The exact binary mixed cubic norm [46]	1593
IV.28	Restriction, exact agreement, and stability [46]	1595
IV.29	Constructive integration and optimal phase order [46]	1598
IV.30	An all-characteristic replacement of scalar polarization [46]	1601
IV.31	Complete classification of trace-multiplication tensors [46]	1602
IV.32	Exact restriction dimensions and repair costs [46]	1604
IV.33	Returning to Gowers' derivative-spectrum step [46]	1606
IV.34	Cube identities and phase invariance [43]	1608
IV.35	Division-free finite-abelian orthogonality [43]	1609
IV.36	Sharp symmetry thresholds on finite abelian groups [43]	1611
IV.37	Boolean nonclassical integration, explicitly [43]	1613
IV.38	The shear identity and sharp Boolean certification [43]	1615
IV.39	Exact cubic obstruction ranks and optimal repair [43]	1617
IV.40	Mixed correlations: an optimal fourth power and a smaller repair bound [43]	1620
IV.41	The quartic frontier: an exact small quotient and a global reduction [43]	1622
IV.42	Source 47's chapter: exact Boolean obstruction energies [write]	1626
IV.43	Finite calculus and nonclassical integration [47]	1626
IV.44	The spectral obstruction supplied by the defect [47]	1628
IV.45	Derivative periods and forbidden frequencies [47]	1630
IV.46	Exact canonical energies in every degree [47]	1631
IV.47	The geometry of low-complexity defects [47]	1633
IV.48	Support bounds for non-pure symmetric multilinear maps [47]	1635
IV.49	Universal thresholds and exact phase removal [47]	1637
IV.50	Deficits, noisy selectors, and retained constraints [47]	1638
IV.51	Sources 65 and 69: Boolean phase integration in every degree [write]	1640
IV.52	Finite calculus and the integration criterion [65]	1641
IV.53	The cubic spectral base case [65]	1643
IV.54	The canonical energy in every degree [65]	1644
IV.55	Defect geometry and the integrable-slice identity [65]	1646
IV.56	Closing the induction in every degree [65]	1648
IV.57	Rigidity, repair geometry, and exact counting [65]	1649
IV.58	Robust certification and retained contraction profiles [65]	1651
IV.59	Finite calculus and the integrable tensor space [69]	1653
IV.60	The defect and the cubic spectral base case [69]	1655

IV.61	Rank-one defects and canonical energies [69]	1657
IV.62	The contraction-integrability induction [69]	1659
IV.63	Extremal flags and quantitative tensor rigidity [69]	1660
IV.64	The minimal obstruction quotient [69]	1662
IV.65	A construction attaining every dimension bound [69]	1664
IV.66	Energy-to-complexity and robust selector certificates [69]	1666
IV.67	Source 63’s Section 6: the face cover at Proposition 17.7 [write]	1667
IV.68	Asymptotically sharp local phase removal [63]	1668

V Lean interface, ledger crosswalk and research questions 1679

V.1	The formal project: ledger crosswalk [write]	1683
V.2	Further questions and research [write]	1707
V.3	Relationship to the ProveIt formalization [01]	1753
V.4	Further research questions and concrete next steps [01]	1755
V.5	Conclusions and result-status ledger [01]	1757
V.6	Boundary and normalization checklist [01]	1758
V.7	Source and artifact provenance [01]	1759
V.8	Formalization interface and reproducibility [02]	1759
V.9	Further research questions and proposed work [02]	1760
V.10	Conclusion [02]	1762
V.11	Source audit and normalization checklist [02]	1762
V.12	How local improvements affect an iteration [03]	1763
V.13	The Lean development and a concrete formalization route [03]	1765
V.14	Further research questions and topics [03]	1768
V.15	Reproducibility and verification [03]	1770
V.16	Quantitative comparisons and what they enable [04]	1771
V.17	Further research questions and a proposed program [04]	1772
V.18	Closing perspective [04]	1776
V.19	A map to the Lean development [04]	1776
V.20	Reproducibility and supplementary checks [04]	1777
V.21	Relation to the Lean development [05]	1779
V.22	Reproducible computation and the limits of testing [05]	1780
V.23	Further research questions [05]	1780
V.24	Conclusion [05]	1782
V.25	An explicit potential for density-dependent iteration [06]	1782
V.26	Compatibility with the Lean development [06]	1785
V.27	Further research questions and a concrete program [06]	1787
V.28	Verification and the Lean integration boundary [07]	1789
V.29	Further research questions and priorities [07]	1791
V.30	Conclusion [07]	1792
V.31	Reproduction commands [07]	1793
V.32	Integration with the supplied Lean development [08]	1793
V.33	Exact checks and reproducibility [08]	1795
V.34	Further research questions [08]	1795
V.35	Conclusion [08]	1798
V.36	Build and source provenance [08]	1798
V.37	Reproducible exact checks [09]	1799
V.38	Integration with the supplied Lean development [09]	1800
V.39	Further research questions and a prioritized program [09]	1802
V.40	Conclusion [09]	1805
V.41	Proof and artifact dependency map [09]	1805
V.42	What the local gains do to a quantitative iteration [10]	1806
V.43	A map to the Lean development and reproducible checks [10]	1807

V.44	Further research questions and proposed next steps [10]	1809
V.45	External inputs, provenance, and reproducibility [10]	1811
V.46	Finite verification and formalization interfaces [11]	1812
V.47	Further research questions [11]	1814
V.48	Conclusion [11]	1815
V.49	Reproduction and provenance [11]	1815
V.50	Formalization interfaces and verification [12]	1816
V.51	Further research questions and topics [12]	1817
V.52	Parameter ledger and reproducibility [12]	1820
V.53	Finite examples, verification, and formalization [13]	1820
V.54	Further research questions and concrete next steps [13]	1823
V.55	Reproduction and source provenance [13]	1826
V.56	The Lean source map and a concrete formalization route [14]	1826
V.57	A focused research program [14]	1829
V.58	Reproducibility, diagnostics, and proof status [14]	1830
V.59	Reproducible computations and Lean integration [15]	1832
V.60	Relation to existing work and remaining research [15]	1834
V.61	Conclusion [15]	1836
V.62	Notation and verification inventory [15]	1837
V.63	Source pin and proof-status boundaries [15]	1838
V.64	How the refinements enter the larger argument [16]	1838
V.65	Source fidelity and a formalization plan [16]	1839
V.66	Reproducible computations and their role [16]	1840
V.67	Further research questions and priorities [16]	1841
V.68	Conventions for reproducibility and citation [16]	1842
V.69	Relationship with the inspected Lean development [17]	1843
V.70	Finite audits and reproducibility [17]	1844
V.71	Further research questions and proposed directions [17]	1845
V.72	Conclusion [17]	1847
V.73	A compact proof-dependency map [17]	1847
V.74	Source provenance and claim ledger [17]	1847
V.75	Quantitative consequences and limits of substitution [18]	1848
V.76	Relation to the Lean development [18]	1850
V.77	Exact finite verification and reproducibility [18]	1852
V.78	Further research questions and proposed next steps [18]	1852
V.79	Source record and attribution [18]	1854
V.80	Source audit and the boundary of the Lean integration [19]	1854
V.81	Verification and reproducibility [19]	1858
V.82	Further research questions and proposed projects [19]	1859
V.83	Formalization and reproducibility [20]	1861
V.84	Further research questions and topics [20]	1863
V.85	Conclusion [20]	1864
V.86	Source audit and package scope [20]	1865
V.87	Repository interface and verification [21]	1865
V.88	Further research questions [21]	1867
V.89	Finite incidence certificates and reproducibility [21]	1869
V.90	Source manifest and status of claims [21]	1870
V.91	Verification, attribution, and repository integration [22]	1871
V.92	Further research questions and proposed next steps [22]	1873
V.93	Verification record and repository integration [23]	1874
V.94	Further research questions and topics [23]	1876
V.95	Build instructions and package contents [23]	1878
V.96	Examples, boundary cases, and reproducible checks [24]	1878
V.97	Integration with the inspected Lean development [24]	1880

V.98	Further research questions and concrete next steps [24]	1882
V.99	Conclusion [24]	1884
V.100	A compact proof and constant ledger [24]	1884
V.101	Package contents and rebuilding [24]	1885
V.102	Repository integration and verification record [25]	1885
V.103	Further research questions [25]	1887
V.104	Conclusion [25]	1888
V.105	Source comparison and remaining questions [26]	1888
V.106	The exact companion and reproducibility [26]	1890
V.107	Verification and repository integration [27]	1890
V.108	Further research questions [27]	1892
V.109	Conclusion [27]	1895
V.110	Reproduction checklist and file inventory [27]	1895
V.111	An explicit interface for further development [28]	1895
V.112	Exact computations and the scope of verification [28]	1897
V.113	Further research questions and topics [28]	1899
V.114	Source provenance and attribution [28]	1902
V.115	How to use these results in the ProveIt project [29]	1903
V.116	Verification and reproducibility [29]	1905
V.117	Further research questions [29]	1906
V.118	Repository crosswalk and exact validation [30]	1910
V.119	Research questions and next integration targets [30]	1912
V.120	Conclusion [30]	1913
V.121	Build, verification, and file layout [30]	1913
V.122	Verification and a proposed formal integration [31]	1914
V.123	Further research questions [31]	1915
V.124	A short dictionary for the extremal geometry [31]	1916
V.125	The pinned source and the literature [32]	1917
V.126	Reproducibility and further questions [32]	1919
V.127	Further research questions and a concrete program [33]	1920
V.128	Repository integration and formalization interface [33]	1922
V.129	Normalization and audit checklist [33]	1924
V.130	Consequences, limitations, and a research program [34]	1925
V.131	Reproducibility and proposed ProveIt integration [34]	1928
V.132	Evidence, source scope, and formalization [35]	1930
V.133	A proposed research programme [35]	1933
V.134	Computational checks and reproducibility [36]	1936
V.135	ProveIt integration and formalization plan [36]	1937
V.136	Further research questions [36]	1939
V.137	Conclusion [36]	1941
V.138	Verification and repository integration [37]	1941
V.139	Further research problems [37]	1943
V.140	Conclusion [37]	1945
V.141	Source audit and reproducibility notes [37]	1945
V.142	What the calculations verify [38]	1946
V.143	Integration with the ProveIt development [38]	1947
V.144	Proposed further research questions and topics [38]	1949
V.145	A precise interface with the inspected Gowers development [39]	1952
V.146	Exact checks, examples, and reproducibility [39]	1955
V.147	Further research questions and proposed next steps [39]	1956
V.148	Conclusion [39]	1958
V.149	Proof dependencies and formalization targets [39]	1958
V.150	Notation and endpoint checklist [39]	1959
V.151	Reproducibility, implementation, and formalization [40]	1960

V.152	Further research questions and integration priorities [40]	1961
V.153	Conclusion [40]	1963
V.154	The pinned source and classical context [41]	1963
V.155	Reproducibility and further questions [41]	1964
V.156	Relation to prior arguments and further questions [42]	1965
V.157	Reproducible bounded diagnostics [42]	1966
V.158	Relation to Gowers’s argument and the formal project [43]	1967
V.159	Exact checks, provenance, and limitations [43]	1968
V.160	Further research questions and topics [43]	1970
V.161	Conclusion [43]	1971
V.162	Verification, source comparison, and formalization targets [44]	1971
V.163	Further research questions and proposed next targets [44]	1973
V.164	Conclusion [44]	1974
V.165	Reproduction and package contents [44]	1975
V.166	Further research questions and a proposed programme [45]	1975
V.167	Normalization and a formalization roadmap for ProveIt [45]	1979
V.168	Exact certificates and reproducibility [45]	1980
V.169	Source ledger and limits of the audit [45]	1982
V.170	Algorithms, reproducibility, and proof status [46]	1983
V.171	Further research questions and concrete next targets [46]	1985
V.172	Conclusion [46]	1986
V.173	Proof dependency map and scope checks [46]	1987
V.174	Relation to the Gowers argument and the formal project [47]	1987
V.175	Exact verification and source provenance [47]	1988
V.176	Further research questions and topics [47]	1989
V.177	Conclusion [47]	1991
V.178	Reproduction and package layout [47]	1991
V.179	Worked instances and reproducible exact checks [48]	1991
V.180	Integration into the Ramsey development [48]	1993
V.181	Further research questions [48]	1995
V.182	Conclusion [48]	1996
V.183	Algorithmic skeletons and invariants [48]	1996
V.184	Proof-dependency map and audit boundaries [48]	1997
V.185	Mathematical consequences and integration boundaries [49]	1998
V.186	Further research questions and a proposed program [49]	1999
V.187	Reproducibility and repository integration [49]	2002
V.188	Computational checks and reproducibility [50]	2003
V.189	A formalization route [50]	2005
V.190	Further research questions [50]	2006
V.191	Conclusion [50]	2007
V.192	Source-to-witness crosswalk [50]	2007
V.193	Definition-level pseudocode for the witness [50]	2008
V.194	Proof and provenance checklist [50]	2008
V.195	The two interfaces after complete upstream extraction [51]	2009
V.196	Proof and computation have different roles [51]	2010
V.197	Pinned source map and comparison scope [51]	2010
V.198	Repairs and questions left open [52]	2012
V.199	Reproducibility and verification limits [52]	2013
V.200	Pinned source map [52]	2015
V.201	A concise reconstruction guide [52]	2015
V.202	Scope and further questions [53]	2016
V.203	What the reproducible companion verifies [53]	2017
V.204	Pinned source map [53]	2017
V.205	Repository integration and proof dependencies [54]	2019

V.206	Further research questions [54]	2020
V.207	Conclusion [54]	2021
V.208	Further research questions [55]	2022
V.209	Validation, provenance, and integration [55]	2023
V.210	A compact coefficient ledger [55]	2025
V.211	Why a purely formal Taylor expansion is insufficient [55]	2026
V.212	Deterministic reconstruction and exact diagnostics [56]	2026
V.213	The Gowers interface and a formalization route [56]	2028
V.214	Further research questions [56]	2031
V.215	Proof audit and dependency ledger [56]	2032
V.216	Further research questions and a proposed programme [57]	2034
V.217	Proof dependencies, verification, and integration [57]	2037
V.218	Examples, numerical exploration, and exact checks [58]	2040
V.219	Further research questions and a proposed programme [58]	2042
V.220	The essential integer certificate [58]	2044
V.221	Package contents and reproduction [58]	2045
V.222	Reproducibility and verification scope [59]	2045
V.223	Further research and a formalization programme [59]	2047
V.224	Consequences, comparison, and integration [60]	2051
V.225	Computational evidence and reproducibility [60]	2053
V.226	Further research questions and a proposed programme [60]	2055
V.227	The inherited coarse decoder, with proof [60]	2058
V.228	Exact computational checks and constructive content [61]	2061
V.229	Integration and proof-assistant targets [61]	2063
V.230	Further research questions [61]	2065
V.231	Conclusion [61]	2066
V.232	Source and validation ledger [61]	2067
V.233	Exact checks and reproducibility [62]	2067
V.234	Repository integration and proof obligations [62]	2068
V.235	Further research questions [62]	2069
V.236	Conclusion [62]	2070
V.237	Further research questions and a proposed program [63]	2070
V.238	Source audit, verification, and integration [63]	2074
V.239	Exact verification, implementation, and repository integration [64]	2076
V.240	Further research questions [64]	2078
V.241	Conclusion [64]	2079
V.242	Explicit certificates and finite verification [65]	2079
V.243	Relation to Gowers’s argument and the formal interfaces [65]	2081
V.244	Further research questions [65]	2082
V.245	Conclusion [65]	2083
V.246	A small explicit certificate [65]	2084
V.247	Reproduction [65]	2084
V.248	Exact recovery algorithms and reproducibility [66]	2084
V.249	Further research questions [66]	2086
V.250	Conclusions [66]	2088
V.251	The relation of the first gap to inverse Pollard theory [66]	2088
V.252	Proof dependencies and formalization plan [66]	2089
V.253	Source audit and reproduction notes [66]	2090
V.254	Further research: precise questions and plausible routes [67]	2091
V.255	Verification, reproducibility, and repository integration [67]	2095
V.256	Consequences, interfaces, and further research [68]	2097
V.257	Verification, provenance, and integration notes [68]	2102
V.258	Finite certificates and implementation interfaces [69]	2105
V.259	Relation to Gowers’s argument and the ProveIt formal project [69]	2106

V.260 Exact verification and source provenance [69]	2107
V.261 Further research questions and topics [69]	2109
V.262 Conclusion [69]	2110
V.263 Reproduction and package layout [69]	2110
References	2112

F.1 Guide to this report

This report merges five research manuscripts, numbered 01–05 in arrival order and then in ASCII order of their archive names. All five refine local estimates of W. T. Gowers’s proof of Szemerédi’s theorem [1], and all five compare their statements with ProveIt’s edited transcription of the paper (`Combinatorics/Ramsey/Papers/sz-thm-gowers-proof`) and its Lean development (`Combinatorics/Ramsey/Lean/GowersSzemerédi`). The report lives beside both, in `Combinatorics/Ramsey/Research/GowersSzemerédi/local-quantitative-refinements`, by Vladimir’s direction; that placement confers no formal status (Section F.5).

No.	Delivered title	Pin	Parts
01	Sharp Local Refinements of Gowers’s Szemerédi Argument: density transfer, carry-pattern entropy, and formalization interfaces	3a25020a6	I, IV, V
02	Energy-Sensitive Inverse Steps and Sharp Density Extraction	66c02f8a1*	I, II, V
03	Sharper local estimates in Gowers’s proof of Szemerédi’s theorem: density increments, exact cube expansions, and dependent random choice	3a25020a6	I, II, III, V
04	Sharper Local Bounds in Gowers’s Proof of Szemerédi’s Theorem: polynomial homomorphism restriction, sharp cube estimates, and floor-pattern complexity (the base)	327773149	II, III, IV, V
05	Sharp Finite Phase Partitions and Density-Sensitive Increments	327773149	I, V
<i>Added in the second write, 6 October 2026</i>			
06	Sharp Local Bounds in Gowers’ Proof of Szemerédi’s Theorem: carry patterns, quadratic counting errors, and optimal density extraction	7d99e5dee	I, III, IV, V
07	Transport-sensitive rigidity in Gowers’s approximate-homomorphism argument	ae0e89574	II, V
08	Fourth-order stability of additive cube counts	none [†]	III, V
09	Arithmetic Circuit Expansions for Cubical Counts	7d99e5dee	III, V
10	Quantitative Refinements Around Gowers’ Proof of Szemerédi’s Theorem: torsion, density profiles, additive energy, and polynomial colorings	7d99e5dee	I, II, III, V
<i>Added in the third write, 6 October 2026</i>			
11	Sharp Affine Rigidity from Graph Additive Energy: cubic stability, target torsion, and exact non-affine extremizers	40f885098	II, V
12	Sharp Freiman Extension and Polynomial Bohr Radii	9ea383e9b	II, V
13	Sharp Local Refinements in Gowers’ Szemerédi Framework: simultaneous recurrence, extremal density selection, and centered cube counts	9ea383e9b	I, III, V
14	Sharp Extremal Refinements in Gowers’s Szemerédi Framework: secondary cube asymptotics, arrangement stability, and polynomial supersaturation	9ea383e9b	II, III, V

No.	Delivered title	Pin	Parts
<i>Added in the fourth write, 6 October 2026</i>			
15	Collision-Sensitive Amplification of Respected Arrangements: sharp refinements of Gowers’s Lemma 14.7, coupled product-property bounds, and Cartesian stability	9ea383e9b	II, V
16	Local Phase Removal Without Carry Patterns, and further quantitative refinements of Gowers’ proof of Szemerédi’s theorem	dca9f0638	I, III, IV, V
17	Polynomial-loss restriction in Gowers’s Szemerédi argument: Fejér filters, density-sensitive moments, and finite-field systems	dca9f0638	II, V
18	Higher-Order Cube Expansions and Sharp Degeneracy Bounds in Gowers’ Szemerédi Framework: simultaneous polynomial partitions, exact constants, and quantitative obstructions	dca9f0638	I, II, III, V
19	Efficient Interpolation in Gowers’ Szemerédi Framework: exact sampling losses, shared refinements, and simultaneous polynomial recurrence	dca9f0638	I, II, V
20	Quadratic Counting Errors and Sharp Directional Bounds in Gowers’s Szemerédi Argument: two-function uniformity, torsion-sensitive constants, and a sharp three-term obstruction	1bd57e91c [‡]	III, V
No.	Delivered title	Pin	Parts
<i>Added in the fifth write, 6 October 2026</i>			
21	Sharp Characteristic-Five Cube Bounds in Gowers’s Framework: optimal sixth-order correction, phase rigidity, and an eighth-order extremal term	43306653a	III, V
22	Degree Filtrations in Gowers’ Ramsey Framework: sharper degeneracy bounds, exact polynomial pattern thresholds, and obstructions to extending norm constructions	43306653a	II, III, V
23	Sharp Thresholds and Extremal Constants in Gowers’s Szemerédi Framework: numerical semigroups, characteristic-five cube asymptotics, and common partitions for affine fibres	65f4e039f	I, II, III, V
24	Sharper local structure in Gowers’s Szemerédi argument: moment-sensitive dependent random choice, small difference sets, and vertical-fibre compression	2e0553b0c	II, V
25	The Codimension-One Degeneracies of Gowers Arrangements: an asymptotically sharp replacement for Lemma 15.4 and explicit random-selection thresholds	568d1802e	II, V
26	Threshold free affine density increments on cyclic groups (Report275)	c94998a5d	I, V
27	Sharp Degeneracy Bounds for Gowers Arrangements: hyperplane obstructions, second-order structure, and an exact finite-field count	none [§]	II, V

No.	Delivered title	Pin	Parts
<i>Added in the sixth write, 6 October 2026</i>			
28	Logarithmic Selection, Local Freiman Models, and Higher Uniformity	e3125e698	II, III, V
29	Further Quantitative Refinements of Gowers's Szemerédi Argument: higher norm comparisons, joint quadratic rigidity, and restriction by affine rank	e3125e698	II, III, V
30	Fixed-radius Bohr invariance in Gowers's Szemerédi argument: escape-tube bounds, larger Freiman-extension neighborhoods, and exponential-rank obstructions	8d65d3c9f	II, V
31	Weighted Derivative Spectra and Sharp Local Rigidity: explicit Freiman restriction, an optimal near-extremal U^3 profile, and the limits of stability for correlation weights	e3125e698	II, V
32	Threshold free quadratic density increments on cyclic groups (Report276)	c94998a5d	I, V
33	Sharp High-Agreement Stability for Gowers's Respected Arrangements: full-product decoding, optimal error terms, and erasures	9f83dbe0f	II, V
No.	Delivered title	Pin	Parts
<i>Added in the seventh write, 6 October 2026</i>			
34	Prime Torsion and Norm Gaps in Gowers Cube Counts: minimal Boolean supports, the cost of positivity, and sharper uniformity inequalities	12da35673	II, III, V
35	Graded Fourier Selection and Two Further Refinements of Gowers' Argument: reduced field-size thresholds, a stronger real U^4 gap, and exact polynomial partition obstructions	12da35673	I, II, III, V
36	Sharp Symmetry Certification for Gowers Derivative Spectra: alternating-rank bounds, lossless polarization, and robust finite-field inverse steps	6c2e2172b	IV, V
37	Sharp Spectral Detection of Missing Intervals: exact moment tradeoffs, extremal rigidity, and quantitative recurrence in Gowers's argument	869289aa3	I, V
38	Exact Quotient Reduction and Sharp Stability in Gowers' Ramsey Framework: characteristic-five cube extremizers, odd-order rigidity, and six-term norm-extension defects	12da35673	II, III, V
39	Optimal Shared-Anchor Interpolation and Labelled Slice Refinement in Gowers's Szemerédi Argument: exact finite minimax bounds, sharp Poisson limits, and a replacement for the sampling step of Lemma 16.10	869289aa3	II, V

No.	Delivered title	Pin	Parts
<i>Added in the eighth write, 7 October 2026</i>			
40	One-stage Fejér selection for Gowers arrangements: single-monomial relation isolation, polynomial parameter loss, and sharp exponents within a positive-Fourier selection class	none	II, V
41	Threshold free polynomial density increments on cyclic groups (Report277)	c94998a5d	I, V
42	Sharp quadratic scale density increments on cyclic groups (Report278)	b16ce380d	I, V
43	Sharp Boolean Phase Integration from Gowers Derivative Spectra: exact cubic obstruction ranks, repeated-variable certification, finite-abelian orthogonality, and a quartic reduction	797aa3cfe	IV, V
44	Finite-Grid Formulae and Exact Core-Avoidance Laws for Gowers Arrangements: computing the second degeneracy coefficient in polynomially many arithmetic operations	8c12c906e	II, V
45	Rigidity gaps in Gowers's Szemerédi framework: exact additive-energy levels, punctured cosets, and a real fourth-uniformity inequality	6ce33539e	II, III, V
46	Sharp Cubical Correlation and Characteristic Barriers in Gowers' Polarization Step: exact binary cubic norms, improved restriction bounds, optimal phase depth, and trace-multiplication tensors	none (blobs)	IV, V
No.	Delivered title	Pin	Parts
<i>Added in the ninth write, 7 October 2026</i>			
47	Exact Boolean Obstruction Energies: derivative periods, sharp quartic–sextic integration, and dimension-free higher-degree extremizers	5c9a442d2	IV, V
48	Collision-aware derandomization of Gowers's restriction step: exact energy inventories, prescribed-cardinality subsets, and constructive restriction for linear systems	0b4890a1c	II, V
49	Second-Order Degeneracy, Dual Uniformity Bounds, and an Acyclic Model Theorem: new refinements in Gowers' Szemerédi framework	8c12c906e	II, III, V
50	Sharp Branch-Cover Profiles and a Counterexample to the Two-Input Form of Lemma 16.10	588552b79	II, V
51	An obstruction and a repair for Section 13 square geometry (Report279)	aad4789fc [†]	II, V
52	Composite modulus obstructions to Stage 13.7 (Report280)	5c9a442d2	II, V
53	Explicit uniform density square extraction with a sharper side length (Report281)	bf2fe146d	II, V
[†] and five later pins up to 4f8cd1198			

No.	Delivered title	Pin	Parts
<i>Added in the tenth write, 7 October 2026</i>			
54	A Single Quadratic Phase Across All Translates: spectral extraction, offset rigidity, and sharp stability in Gowers's Section 8	6a6d7961c	II, V
55	Sharp Second-Order Stability at the Polynomial-Phase Endpoint of Gowers Norms: a parity-sensitive refinement of the Szemerédi framework	128f514af	II, V
56	Nearly Sharp Relative Affine Repair on Fourier-Uniform Domains: weighted Freiman tests, error-independent uniformity, and explicit quadratic-level-set consequences	b6a8ee966	II, V
57	Quantitative Bilinear Extraction and Bounded Additive Defects: refinements of Gowers's argument and the ProveIt Ramsey research programme	128f514af	II, V
58	Polynomial restriction on every cyclic group: a torsion-sensitive refinement of Gowers's Lemma 9.3 and an optimal exponent for positive Fourier filtering	5d4718772	II, V
59	Sharp Character Separation and Affine Rigidity in Gowers' Inverse Steps: torsion-sensitive Freiman restrictions, signed energy inequalities, and a research programme for the ProveIt Ramsey project	bec460a33	II, V

No.	Delivered title	Pin	Parts
<i>Added in the eleventh write, 7 October 2026</i>			
60	Fourier limits, exact arrangement repair, and translated restriction in Gowers’s quantitative framework	89c1cbc38	II, III, V
61	Exact Freiman Restriction in Arbitrary Finite Abelian Groups: torsion-uniform coloring, relation-energy retention, and logarithmic order barriers	5ea1d06dd	II, V
62	Unit-Dilation Restriction on Cyclic Groups: polynomial losses for Gowers’ approximate homomorphisms and exact weighted degeneracy constants	17123532a	II, V
63	Asymptotically Sharp Constants in Gowers’ Szemerédi Framework: odd-order norm limits and efficient local phase removal	d7b83fa1f	III, IV, V
64	Binomial separation and endpoint bounds for uniform progression counts: a six-variable digit construction, an exact nine-copy transfer, and obstructions to perfect cyclic codes	none [†]	III, V
65	Sharp Boolean Phase Integration in Every Degree: defect-radical slicing, exact nonintegrability thresholds, and extremal tensor rigidity in Gowers’s framework	8c40adc24	IV, V
66	Partition Normal Forms for Small Schur Defect: sharp endpoint bounds and the first four additive-energy levels	none [†]	II, V
67	Exact Local Profiles and Structural Barriers in Gowers’ Szemerédi Framework: unequal Fourier budgets, characteristic-three rigidity, optimal interpolation lists, and norm layers without a degree cap	none [†]	I, II, III, V
68	Sharp energy gaps and efficient restriction: coset rigidity, puncture stability, and Fourier-filter bounds in Gowers’ Szemerédi framework	0d93c1eac	II, V
69	Sharp All-Degree Boolean Phase Integration: contraction radicals, extremal flags, and optimal finite-dimensional obstruction models	0d93c1eac	IV, V

* source 02 records no pin; this is the branch reference it observed (its Appendix A). [†] sources 64, 66 and 67 record no pin; they identify what they read by blob and SHA-256 (Section F.4, eleventh write). [†] source 08 records blob identifiers only; they are those of 7d99e5dee. [‡] source 20 records an observed branch checkpoint (its Appendix A), not a pin. [§] source 27 records blob identifiers only (its SOURCE_AUDIT.md); all three are unchanged at HEAD.

Arrangement. Every manuscript spans two to four themes, so the report is arranged by theme, not by manuscript. Part I treats density transfer and phase-flat partitions; Part II the inverse step (energies, moments, restriction, dependent random choice); Part III cube and progression counts; Part IV floor patterns and the phase extraction of Proposition 17.7; Part V the formal interface, the crosswalk to the ledger, and the research questions. Inside a Part the sources follow in numerical order. Each source’s sections are printed whole, as delivered, with their own subsections; a section heading carries its source tag (for example “[01]”) and a line naming the delivered section. Each manuscript’s abstract and opening section are printed in the front matter (Sections F.6–F.10), where its title block is quoted.

Numbering. Sections restart in each Part and carry its numeral: Section I.9 is the ninth section

of Part I, and the front matter's sections are F.1–F.10. Statements and equations are numbered within sections, so source 02's Theorem 7.1 is Theorem I.9.1 here. A statement or equation $k.j$ of a source keeps its index j within its section, except that sources 02 and 03 numbered their equations consecutively through the whole manuscript, whereas here they are numbered within sections. Source 04's tags (M1)–(M7), (R1)–(R32), (A1)–(A14), (Q1), (Q2) and (*) are kept. *Bare numbers such as “Lemma 5.15” or “Corollary 2.4” always refer to Gowers’s paper* (in the corrected transcription when the text says so), as in all five sources; this report’s own numbers always carry a Part numeral or an F. Plain-text references of a source to its own sections were made live in one place (source 02, Section II.4); everywhere else the sources refer to their own numbers through cross-references, which now print this report’s numbers.

Concordance. Delivered sections and the sections here:

Source	Delivered section → section here
01	1→F.6, 2→I.1, 3→I.2, 4→I.3, 5→I.4, 6→I.5, 7→I.6, 8→I.7, 9→IV.1, 10→IV.2, 11→IV.3, 12→IV.4, 13→IV.5, 14→V.3, 15→V.4, 16→V.5, App. A→V.6, App. B→V.7
02	1→F.7, 2→II.1, 3→II.2, 4→II.3, 5→II.4, 6→I.8, 7→I.9, 8→I.10, 9→I.11, 10→V.8, 11→V.9, 12→V.10, App. A→V.11, App. B→I.12
03	1→F.8, 2→I.13, 3→I.14, 4→I.15, 5→III.1, 6→III.2, 7→II.5, 8→II.6, 9→V.12, 10→V.13, 11→V.14, App. A→V.15
04	1→F.9, 2→II.7, 3→II.8, 4→II.9, 5→II.10, 6→III.3, 7→III.4, 8→IV.6, 9→V.16, 10→V.17, 11→V.18, App. A→V.19, App. B→V.20
05	1→F.10, 2→I.16, 3→I.17, 4→I.18, 5→I.19, 6→I.20, 7→V.21, 8→V.22, 9→V.23, 10→V.24
06	1→F.11, 2→IV.7, 3→III.5, 4→IV.8, 5→I.21, 6→V.25, 7→IV.9, 8→V.26, 9→V.27
07	1→F.12, 2→II.13, 3→II.14, 4→II.15, 5→II.16, 6→II.17, 7→II.18, 8→II.19, 9→V.28, 10→V.29, 11→V.30, App. A→II.12, App. B→V.31
08	1→F.13, 2→III.6, 3→III.7, 4→III.8, 5→III.9, 6→III.10, 7→III.11, 8→III.12, 9→III.13, 10→III.14, 11→V.32, 12→V.33, 13→V.34, 14→V.35, App. A→III.15, App. B→V.36

Source	Delivered section → section here
09	1→F.14, 2→III.16, 3→III.17, 4→III.18, 5→III.19, 6→III.20, 7→III.21, 8→III.22, 9→III.23, 10→III.24, 11→V.37, 12→V.38, 13→III.25, 14→V.39, 15→V.40, App. A→III.26, App. B→III.27, App. C→V.41
10	1→F.15, 2→I.22, 3→I.23, 4→I.24, 5→III.28, 6→III.29, 7→II.20, 8→II.21, 9→III.30, 10→V.42, 11→V.43, 12→V.44, App. A→V.45
11	1→F.16, 2→II.23, 3→II.24, 4→II.25, 5→II.26, 6→II.27, 7→II.28, 8→II.29, 9→II.30, 10→II.31, 11→II.32, 12→V.46, 13→V.47, 14→V.48, App. A→V.49
12	1→F.17, 2→II.34, 3→II.35, 4→II.36, 5→II.37, 6→II.38, 7→II.39, 8→II.40, 9→II.41, 10→V.50, 11→V.51, App. A→II.42, App. B→V.52
13	1→F.18, 2→I.25, 3→I.26, 4→I.27, 5→I.28, 6→I.29, 7→III.31, 8→III.32, 9→III.33, 10→V.53, 11→V.54, App. A→V.55
14	1→F.19, 2→III.34, 3→III.35, 4→II.43, 5→II.44, 6→III.36, 7→III.37, 8→V.56, 9→V.57, App. A→III.38, App. B→V.58
15	1→F.20, 2→II.46, 3→II.47, 4→II.48, 5→II.49, 6→II.50, 7→II.51, 8→II.52, 9→II.53, 10→II.54, 11→II.55, 12→II.56, 13→V.59, 14→V.60, 15→V.61, App. A→V.62, App. B→V.63
16	1→F.21, 2→IV.11, 3→IV.12, 4→IV.13, 5→IV.14, 6→I.30, 7→I.31, 8→III.39, 9→V.64, 10→V.65, 11→V.66, 12→V.67, App. A→V.68
17	1→F.22, 2→II.58, 3→II.59, 4→II.60, 5→II.61, 6→II.62, 7→II.63, 8→II.64, 9→V.69, 10→V.70, 11→V.71, 12→V.72, App. A→V.73, App. B→V.74
18	1→F.23, 2→II.66, 3→III.40, 4→I.32, 5→I.33, 6→V.75, 7→V.76, 8→V.77, 9→V.78, App. A→V.79
19	1→F.24, 2→II.67, 3→II.68, 4→II.69, 5→II.70, 6→I.34, 7→V.80, 8→II.71, 9→V.81, 10→V.82
20	1→F.25, 2→III.42, 3→III.43, 4→III.44, 5→III.45, 6→III.46, 7→III.47, 8→III.48, 9→V.83, 10→V.84, 11→V.85, App. A→V.86
Source	Delivered section → section here
21	1→F.26, 2→III.50, 3→III.51, 4→III.52, 5→III.53, 6→III.54, 7→III.55, 8→III.56, 9→III.57, 10→V.87, 11→V.88, App. A→V.89, App. B→V.90
22	1→F.27, 2→II.82, 3→II.83, 4→II.84, 5→II.85, 6→III.61, 7→III.62, 8→III.63, 9→III.64, 10→III.65, 11→V.91, 12→V.92
23	1→F.28, 2→I.35, 3→I.36, 4→I.37, 5→I.38, 6→I.39, 7→III.58, 8→II.103, 9→V.93, 10→V.94, App. A→III.59, App. B→V.95
24	1→F.29, 2→II.73, 3→II.74, 4→II.75, 5→II.76, 6→II.77, 7→II.78, 8→II.79, 9→II.80, 10→V.96, 11→V.97, 12→V.98, 13→V.99, App. A→V.100, App. B→V.101
25	1→F.30, 2→II.86, 3→II.87, 4→II.88, 5→II.89, 6→II.90, 7→II.91, 8→II.92, 9→V.102, 10→V.103, 11→V.104
26	1→F.31, 2→I.41, 3→I.42, 4→I.43, 5→V.105, 6→V.106, App. A→I.44
27	1→F.32, 2→II.93, 3→II.94, 4→II.95, 5→II.96, 6→II.97, 7→II.98, 8→II.99, 9→II.100, 10→II.101, 11→II.102, 12→V.107, 13→V.108, 14→V.109, App. A→V.110

Source	Delivered section → section here
28	1→F.33, 2→II.122, 3→II.123, 4→II.124, 5→II.125, 6→II.126, 7→II.127, 8→III.70, 9→V.111, 10→V.112, 11→V.113, 12→V.114
29	1→F.34, 2→III.67, 3→III.68, 4→III.69, 5→II.110, 6→II.111, 7→II.112, 8→II.113, 9→II.120, 10→V.115, 11→V.116, 12→V.117
30	1→F.35, 2→II.128, 3→II.129, 4→II.130, 5→II.131, 6→II.132, 7→II.133, 8→II.134, 9→II.135, 10→V.118, 11→V.119, 12→V.120, App. A→V.121
31	1→F.36, 2→II.105, 3→II.106, 4→II.107, 5→II.108, 6→II.115, 7→II.116, 8→II.117, 9→II.109, 10→V.122, 11→V.123, App. A→II.118, App. B→V.124
32	1→F.37, 2→I.45, 3→I.46, 4→I.47, 5→I.48, 6→V.125, 7→V.126
33	1→F.38, 2→II.137, 3→II.138, 4→II.139, 5→II.140, 6→II.141, 7→II.142, 8→II.143, 9→V.127, 10→V.128, App. A→V.129
Source	Delivered section → section here
34	1→F.39, 2→III.74, 3→III.75, 4→III.76, 5→III.77, 6→III.78, 7→III.79, 8→III.80, 9→III.81, 10→III.71, 11→V.130, 12→V.131, App. A→II.163
35	1→F.40, 2→II.145, 3→II.146, 4→II.147, 5→II.148, 6→III.72, 7→I.61, 8→V.132, 9→V.133
36	1→F.41, 2→IV.16, 3→IV.17, 4→IV.18, 5→IV.19, 6→IV.20, 7→IV.21, 8→IV.22, 9→IV.23, 10→IV.24, 11→V.134, 12→V.135, 13→V.136, 14→V.137
37	1→F.42, 2→I.50, 3→I.51, 4→I.52, 5→I.53, 6→I.54, 7→I.55, 8→I.56, 9→I.57, 10→V.138, 11→V.139, 12→V.140, App. A→I.58, App. B→I.59, App. C→V.141
38	1→F.43, 2→III.83, 3→III.84, 4→III.85, 5→III.86, 6→II.150, 7→III.88, 8→V.142, 9→V.143, 10→V.144
39	1→F.44, 2→II.152, 3→II.153, 4→II.154, 5→II.155, 6→II.156, 7→II.157, 8→II.158, 9→II.159, 10→II.160, 11→II.161, 12→V.145, 13→V.146, 14→V.147, 15→V.148, App. A→V.149, App. B→V.150
Source	Delivered section → section here
40	1→F.45, 2→II.165, 3→II.166, 4→II.167, 5→II.168, 6→II.169, 7→II.170, 8→II.171, 9→II.172, 10→V.151, 11→V.152, 12→V.153
41	1→F.46, 2→I.63, 3→I.64, 4→I.65, 5→I.66, 6→I.67, 7→I.68, 8→V.154, 9→V.155
42	1→F.47, 2→I.69, 3→I.70, 4→I.71, 5→I.72, 6→I.73, 7→I.74, 8→V.156, 9→V.157
43	1→F.48, 2→IV.34, 3→IV.35, 4→IV.36, 5→IV.37, 6→IV.38, 7→IV.39, 8→IV.40, 9→IV.41, 10→V.158, 11→V.159, 12→V.160, 13→V.161
44	1→F.49, 2→II.174, 3→II.175, 4→II.176, 5→II.177, 6→II.178, 7→II.179, 8→II.180, 9→II.181, 10→II.182, 11→V.162, 12→V.163, 13→V.164, App. A→V.165
45	1→F.50, 2→II.184, 3→II.185, 4→III.90, 5→III.91, 6→V.166, App. A→V.167, App. B→V.168, App. C→V.169
46	1→F.51, 2→IV.26, 3→IV.27, 4→IV.28, 5→IV.29, 6→IV.30, 7→IV.31, 8→IV.32, 9→IV.33, 10→V.170, 11→V.171, 12→V.172, App. A→V.173

Source	Delivered section → section here
47	1→F.52, 2→IV.43, 3→IV.44, 4→IV.45, 5→IV.46, 6→IV.47, 7→IV.48, 8→IV.49, 9→IV.50, 10→V.174, 11→V.175, 12→V.176, 13→V.177, App. A→V.178
48	1→F.53, 2→II.187, 3→II.188, 4→II.189, 5→II.190, 6→II.191, 7→II.192, 8→II.193, 9→II.194, 10→V.179, 11→V.180, 12→V.181, 13→V.182, App. A→V.183, App. B→V.184
49	1→F.54, 2→II.221, 3→II.222, 4→II.223, 5→II.224, 6→II.225, 7→III.93, 8→II.235, 9→V.185, 10→V.186, 11→V.187
50	1→F.55, 2→II.227, 3→II.228, 4→II.229, 5→II.230, 6→II.231, 7→II.232, 8→II.233, 9→II.234, 10→V.188, 11→V.189, 12→V.190, 13→V.191, App. A→V.192, App. B→V.193, App. C→V.194
51	1→F.56, 2→II.203, 3→II.204, 4→II.205, 5→II.206, 6→II.207, 7→II.208, 8→II.209, 9→II.210, 10→II.211, 11→V.195, 12→V.196, App. A→V.197
52	1→F.57, 2→II.212, 3→II.213, 4→II.214, 5→II.215, 6→II.216, 7→II.217, 8→II.218, 9→II.219, 10→V.198, 11→V.199, App. A→V.200, App. B→V.201
53	1→F.58, 2→II.196, 3→II.197, 4→II.198, 5→II.199, 6→II.200, 7→II.201, 8→II.202, 9→V.202, 10→V.203, App. A→V.204
Source	Delivered section → section here
54	1→F.59, 2→II.268, 3→II.269, 4→II.270, 5→II.271, 6→II.272, 7→II.273, 8→II.274, 9→II.275, 10→V.205, 11→V.206, 12→V.207
55	1→F.60, 2→II.237, 3→II.238, 4→II.239, 5→II.240, 6→II.241, 7→II.242, 8→II.243, 9→II.244, 10→II.245, 11→II.246, 12→V.208, 13→V.209, App. A→V.210, App. B→V.211
56	1→F.61, 2→II.253, 3→II.254, 4→II.255, 5→II.256, 6→II.257, 7→II.258, 8→II.259, 9→II.260, 10→V.212, 11→V.213, 12→V.214, App. A→V.215
57	1→F.62, 2→II.288, 3→II.289, 4→II.290, 5→II.291, 6→II.292, 7→V.216, 8→V.217
58	1→F.63, 2→II.277, 3→II.278, 4→II.279, 5→II.280, 6→II.281, 7→II.282, 8→II.283, 9→II.284, 10→II.285, 11→II.286, 12→V.218, 13→V.219, App. A→V.220, App. B→V.221
59	1→F.64, 2→II.262, 3→II.263, 4→II.264, 5→II.265, 6→II.266, 7→II.249, 8→II.250, 9→II.251, 10→II.252, 11→V.222, 12→V.223
Source	Delivered section → section here
60	1→F.65, 2→III.99, 3→II.332, 4→II.333, 5→II.334, 6→II.335, 7→II.336, 8→II.337, 9→II.330, 10→V.224, 11→V.225, 12→V.226, App. A→V.227
61	1→F.66, 2→II.322, 3→II.323, 4→II.324, 5→II.325, 6→II.326, 7→II.327, 8→II.328, 9→II.329, 10→V.228, 11→V.229, 12→V.230, 13→V.231, App. A→V.232
62	1→F.67, 2→II.314, 3→II.315, 4→II.316, 5→II.317, 6→II.318, 7→II.319, 8→II.320, 9→II.321, 10→V.233, 11→V.234, 12→V.235, 13→V.236
63	1→F.68, 2→III.95, 3→III.96, 4→III.97, 5→III.98, 6→IV.68, 7→V.237, App. A→V.238
64	1→F.69, 2→III.101, 3→III.102, 4→III.103, 5→III.104, 6→III.105, 7→III.106, 8→III.107, 9→V.239, 10→V.240, 11→V.241
65	1→F.70, 2→IV.52, 3→IV.53, 4→IV.54, 5→IV.55, 6→IV.56, 7→IV.57, 8→IV.58, 9→V.242, 10→V.243, 11→V.244, 12→V.245, App. A→V.246, App. B→V.247
66	1→F.71, 2→II.300, 3→II.301, 4→II.302, 5→II.303, 6→II.304, 7→II.305, 8→II.306, 9→II.307, 10→II.308, 11→II.309, 12→V.248, 13→V.249, 14→V.250, App. A→V.251, App. B→V.252, App. C→V.253
67	1→F.72, 2→I.76, 3→II.294, 4→II.295, 5→II.339, 6→III.109, 7→V.254, 8→V.255
68	1→F.73, 2→II.297, 3→II.298, 4→II.299, 5→II.311, 6→II.312, 7→II.313, 8→V.256, 9→V.257
69	1→F.74, 2→IV.59, 3→IV.60, 4→IV.61, 5→IV.62, 6→IV.63, 7→IV.64, 8→IV.65, 9→IV.66, 10→V.258, 11→V.259, 12→V.260, 13→V.261, 14→V.262, App. A→V.263

Notes and edits. Notes added in the write are set in small type, tagged [write] and dated

6 October 2026. They credit results that several sources prove, mark second routes, correct stale statements about the repository against HEAD, record questions answered by another source, and say what the write checked. The write proves four small statements itself, each inside a note with its proof: the comparison of the three Fourier increments (after Table 3), the identification of the four forms of the one-sided bound (after Theorem I.4.1), the value $c_3^{\text{odd}} = 2^{-1/2}$ (after Theorem III.3.4), and the comparison of the two elementary lower bounds for C_k (after Theorem IV.6.2). No delivered statement or proof was changed or dropped. The write’s interventions in delivered text are all of the following kind:

- labels prefixed (`gsr:part:source:`); a label on every delivered section (`...:sec:dk`) and on two remarks of source 03 (`gsr:dt:03:rem:repo`, `gsr:cc:03:rem:rectify`); citation keys unified into one bibliography;
- section headings tagged with the source; theorem styles unified (remarks of sources 04 and 05 were set in the definition style); the sources’ page layout (title pages, running heads, colours, fonts) not reproduced; source 01’s status paragraph set as a small quotation (listed after the independent check of 6 October 2026);
- source 02’s plain “Section 3” made a live reference (above), the two figure paths renamed to the shipped files, and plain-text bookmark forms (`\texorpdfstring`) given to two headings with mathematics, invisible in print;
- source 04’s missing backslash in `_{\epsilon\in\Omega_k}` restored, twice, with a dated note (Section IV.6).

The final section of Part V that the sources’ own questions feed, Section V.2, applies Vladimir’s standing rule of 4 October 2026.

Five further sources. Five more manuscripts on the same subject (sources 06–10, arrival 62e21161f) were placed in this directory by ad37962c4; their prefixed code, data and figures are present. They are not used in this text. A later write will add them to these Parts as dated additions, without renumbering anything printed here.

[write] 6 October 2026; sources 06–10. *Done in the second write.* Sources 06–10 are now part of the report (rows above). Every delivered section of each of them is printed whole in exactly one place: its abstract and Section 1 in the front matter (Sections F.11–F.15, so the front matter now runs to F.15), the other sections appended at the end of Parts I–V after the sections of sources 01–05, in source order, with source 07 as its own chapter of Part II (Section II.11). Nothing printed for sources 01–05 was renumbered or relabelled; the concordance above has gained the rows 06–10. Each Part head and Sections F.2–F.5 and V.1–V.2 have dated additions; results proved by several sources carry dated credit notes at both places. So this paragraph and the opening sentence of this section (“five research manuscripts”) describe the first write only; the title page now says ten.

Further sources. While this write was in progress, sources 11–14, 15–17 and 18–20 were placed in this directory by 6f530c29c, eed2ab863 and 58aa0c493 respectively. Their prefixed files are present (README). They are not used in this text. The placement of sources 11–14 records source 14 as settling source 04’s Question 1 (item III.3); that finding is for the write that adds it.

Edits of the second write in delivered text of sources 06–10 (in addition to the kinds listed above): source 06’s macro `\snapshot` (its own pin) renamed `\snapshotsix`; source 09’s ‘question’ environment (printed “Research question”) set as this report’s ‘Research question’; source 10’s three damaged subscripts M_{rest} restored (note in Section III.28); the two figure paths renamed to the shipped files; source 07’s Appendix A (its notation table) placed at the start of its chapter. Tables and figures of sources 06–10 are numbered within their sections (“Table III.13.1”), and the write’s remark on the false quartic bound is numbered W1, so that no number printed for sources 01–05 moved. Source 06 numbered its equations consecutively through the manuscript; here, as for sources 02 and 03, they are numbered within sections, while every statement of sources 06–10 keeps its index within its section. Source 10’s invisible fragment comments are kept.

Statements proved in the second write, each inside a dated note with its proof: the comparison of source 01’s increment with source 10’s frontier (after Table 3); source 06’s fixed- M bound against the $M-1$ corollary, strict exactly for $\alpha < 1$ (after Corollary I.9.3); the order of the leading step-count coefficients of sources 06, 10 and 03 (after Proposition V.12.1); the comparison of the four progression-existence criteria (after Corollary III.4.4); and the implications from source 10’s Lemma 7.5 and Corollary 7.6 replacements to the formal statements (after Corollary II.21.2). The counterexample to the quartic lower bound (Remark W1) and source 09’s determinant-five configuration were recomputed independently.

[write] 6 October 2026; sources 11–14. *Done in the third write.* Sources 11–14 (placed by 6f530c29c) are now part of the report (rows above); the paragraph of the second write that calls them “not used in this text” describes that write only, as do this section’s opening (“five research manuscripts”) and the second write’s “ten”; the title page now says fourteen. Every delivered section of sources 11–14 is printed whole in exactly one place: their abstracts and Sections 1 in the front matter (Sections F.16–F.19, so the front matter runs to F.19), the other sections appended at the end of Parts I, II, III and V after everything printed there, in source order. Source 11 is a chapter of Part II (Section II.22); logically it follows source 02’s Proposition 6.1 material (Sections II.1–II.4). Source 12 is a chapter of Part II (Section II.33); logically it continues source 07’s chapter on Lemma 10.9 (Section II.11). Both stand at the end of Part II because no number printed earlier may move; their guides say where they belong. Source 13’s partition, selection and increment sections are in Part I, its cube, interval and telescoping sections in Part III; source 14’s Sections 4–5 (coset stability and arrangement equality) are in Part II and its cube, colouring and example sections in Part III. Nothing printed for sources 01–10 was renumbered or relabelled; the concordance has gained the rows 11–14.

Edits of the third write in delivered text of sources 11–14 (in addition to the kinds listed above): sources 11 and 13’s ‘question’ environments (printed “Research question”) set as this report’s ‘Research question’; source 12’s `\snapshot` renamed `\snapshottwelve`; source 12’s ragged-right `\tabularx` columns set inside its own tables; sources 12 and 13’s `\cleveref` references written out with the names `\cleveref` prints (lower-case “lemma”, “theorem”, “equation”, and plurals for lists), because the installed `\cleveref` cannot name theorems that share a counter; the four figure paths renamed to the shipped files; the title-page statements of sources 13 and 14 printed as quotations after their abstracts. Sources 11, 12 and 13 numbered their equations through the whole manuscript; here, as for sources 02, 03 and 06, they are numbered within sections. Every statement keeps its index within its section; tables and figures are numbered within their sections, as for sources 06–10. The third write’s remarks have their own counter and are numbered W2–W4, so that Remark W1 keeps its number.

Proved at the third write, each with its proof: Remark W2 (W2: the formula of source 11’s Theorem F.16.3 fails for $\mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/2\mathbb{Z}$, so its cutoff 34 cannot simply be dropped); Remark W3 (W3: source 11’s near-extremizer bounds composed with source 14’s Theorem II.43.2); Remark W4 (W4: source 14’s Theorem III.34.1 gives the limit in source 04’s display (Q1)); the domination of source 13’s Fourier increment by source 01’s at the same length (note after Theorem I.29.1). Also checked: the scope of source 11’s Corollary II.32.2 (no primality needed), the unused hypotheses of source 12’s Theorem II.35.2, and the proof-level α^5 fraction in the formal proof of Theorem 10.13 (note in Section F.17).

Still pending. Sources 15–17 (eed2ab863), 18–20 (58aa0c493) and 21–23 (a902613c9, placed while this write was in progress) are in this directory and not used in this text. The placement of sources 21–23 records source 21 as answering source 14’s question on the exponent-five coefficient and source 23 as answering source 13’s Question 1 (item I.12); those findings are for the write that adds them, and the items below still state the questions as open.

[write] 6 October 2026; sources 15–20. *Done in the fourth write.* Sources 15–17 (placed by eed2ab863) and 18–20 (placed by 58aa0c493) are now part of the report (rows above); the third write’s “Still pending” paragraph describes that write only, and the title page now says twenty. Every delivered section of sources 15–20 is printed whole in exactly one place: their abstracts and Sections 1 in the front matter (Sections F.20–F.25, so the front matter runs to F.25), the other sections appended at the end of Parts I–V after everything printed there, in source order. Five chapter guides were written for this report: source 15’s (Section II.45) and source 17’s (Section II.57) in Part II, which logically follow source 04’s Sections II.10 and II.8–II.9; sources 18 and 19 on Sections 15 and 16 of the paper (Section II.65); source 20’s in Part III (Section III.41),

beside source 06’s Section III.5; and source 16’s in Part IV (Section IV.10), which proves the strongest of the four Proposition 17.7 coefficients printed there. All stand at the end of their Parts because no number printed earlier may move. Source 16 is split over four Parts (own-step partitions in Part I, cube coefficients in Part III, Proposition 17.7 in Part IV); sources 18 and 19’s polynomial partitions are in Part I. Nothing printed for sources 01–14 was renumbered or relabelled; the concordance has gained the rows 15–20.

Edits of the fourth write in delivered text of sources 15–20 (in addition to the kinds listed above): six macros renamed where the report’s name has another meaning (16’s `\conj` and `\BB`, 17’s `\code`, 18’s `\T`, 19’s `\commit`, 20’s `\U`; the renamed macros print as delivered); source 19’s ‘question’ environment (“Research question”) set as this report’s ‘Research question’, and source 20’s ‘question’ environment, numbered on its own counter there, numbered within its section here (its Questions 1–12 keep their indices); source 19’s `\cleveref` references written out as `\cleveref` prints them; the three figure paths renamed to the shipped files; the title-page statements of sources 17 to 20 (17’s framed status box, 18’s keywords, the closing lines of 19’s and 20’s title pages) printed in their front-matter sections. Source 16 numbers its equations by explicit tags ($k.j$); they are kept, so its equation numbers print as delivered. Tables and figures are numbered within their sections, as for sources 06–14. One write remark was not needed; the counter would continue at W5.

Proved at the fourth write, each inside a dated note with its proof: source 16’s coefficient exceeds the formal $2^{-2(k+1)^3}$ for every k , so its theorem implies `proposition_17_7` (note in Section F.21); the conversion of source 04’s Corollary IV.6.8 to the normalization of the formal statement (91.07 bits at $k = 8$ in place of the 121.66 compared earlier; note after that corollary); source 17’s theorem implies the prime case of `lemma_9_3` and `corollary_9_4` (note in Section F.22), and its constant beats source 04’s exactly for $d \leq 97$ (note after Theorem II.61.2); source 18’s finite bound lies below the formal bound of `lemma_15_4` for $d \geq 2$, $p > 2d$ (note after Theorem II.66.8); source 15’s coupled exponents beat `lemma_14_8` and source 04’s (A12) for every $k \geq 1$ (note after Theorem II.49.6); source 14’s sharp stability implies source 15’s three constants (note after Theorem II.52.2); source 20’s coefficient is strictly below source 06’s (note after Corollary III.45.4). Also checked: source 18’s prescribed-length exponent against source 19’s at the same input (269 against 314 at $q = 1$), the Stirling form of source 18’s Q_d and S_d , and every Lean file and line cited by sources 15–20 and by this write.

Still pending. Sources 21–23 (`a902613c9`), 24–27 (`79d9075b9`) and 28–33 (`c5513046e`; the last two placed while this write was in progress) are in this directory and not used in this text. Their placements record, among other findings, source 21 as answering source 14’s exponent-five questions, source 22 as treating source 18’s $d = 1$ case, source 23 as answering source 13’s Question 1, sources 25 and 27 as re-deriving source 18’s leading degeneracy coefficient, with source 27 answering its second-order question for $d = 2$, source 29 as bounding the odd-order constants c_d^{odd} and adding a restriction theorem beside source 04’s and source 17’s, and source 33 as continuing the arrangement material of sources 14 and 15; those findings are for the writes that add them (Section V.2).

[write] **6 October 2026; sources 21–27.** *Done in the fifth write.* Sources 21–23 (placed by `a902613c9`) and 24–27 (placed by `79d9075b9`) are now part of the report (rows above); the fourth write’s “Still pending” paragraph describes that write only, and the title page now says twenty-seven. Every delivered section of sources 21–27 is printed whole in exactly one place: their abstracts and Sections 1 in the front matter (Sections F.26–F.32, so the front matter runs to F.32), the other sections appended at the end of Parts I, II, III and V after everything printed there. Five chapter guides were written for this report. In Part I, the chapter “Corollary 5.8 without a scale threshold” (Section I.40) holds source 26 and is planned for more sources: its degree-two sibling, source 32 (placed by `c5513046e`, write pending), belongs directly after it, and sources 41 and 42 (placed by `2dfe46f44`, write pending) join the chapter after source 32. In Part II, source 24’s chapter on Section 7 (Section II.72), which logically follows source 03’s Section II.5 and source 10’s Sections II.20–II.21, and a guide to sources 22, 25 and 27 on Lemma 15.4 and source 23 on Lemma 16.10 (Section II.81), which continue source 18’s Section II.66 and source 19’s chapter. In Part III, source 21’s chapter on the exponent-five coefficient (Section III.49), with source 23’s independent proof of the same theorem as a credited second route, which logically follows source 14’s Section III.35; and source 22’s chapter on polynomial colourings (Section III.60), which continues source 10’s Section III.30 and source 14’s Section III.36. Source 23’s partition sections stand at the end of Part I; logically they follow source 13’s Section I.26. All stand at the end of their Parts because no number printed earlier may move. Nothing printed for sources 01–20 was renumbered or

relabelled; the concordance has gained the rows 21–27 (it is now set as three tabulars, rows 01–08, 09–20 and 21–27, so that it can break across pages).

Edits of the fifth write in delivered text of sources 21–27 (in addition to the kinds listed above): three macros renamed where the report’s name has another meaning (23’s `\snapshot`, 24’s `\code`, 25’s `\card`; they print as delivered), and unused macros of 21, 22, 25 and 27 not defined; the ‘question’ environments of sources 22 and 27 (“Research question”) set as this report’s ‘Research question’, those of sources 21, 24 and 25 (“Question”) as its ‘Question’, source 25’s numbered within its section (its own counter there; its Questions 1–8 keep their indices); source 27’s `\cleveref` references written out as its own `\crefname` settings print them (capitalized names; equations as a parenthesized number); source 27’s label `sec:d2` (its Section 7), which has the form of this report’s section labels, prefixed as `gsr:en:27:sec:dtwo`; four labels added to statements delivered without one (source 22’s two-cube theorem and two corollaries of its Section 4, source 23’s indicator proposition), so that dated notes can cite them; the three figure paths renamed to the shipped files; the title-page statements of sources 21, 23, 24 and 27 printed in their front-matter sections. Sources 22 and 23 tag some equations explicitly ((D1.1)–(D1.7); (F1), ...); the tags are kept. Tables and figures are numbered within their sections, as for sources 06–20. No write remark was needed; the counter would continue at W5.

Proved or checked at the fifth write, each inside a dated note with its argument: the degree-one instance of the open `corollary_5_8` holds trivially (by a singleton) whenever $N/M \leq 2^{48}$, and source 26’s Corollary I.41.3 implies that instance in general (note in Section F.31); source 27’s second coefficient $(k^2 + 21k - 6)/2$ from its general formula and $b_{2,2,p} = 3$ (note after Theorem F.32.3); at Lemma 15.5 source 22’s thresholds have the leading term of source 18’s corollary or a worse one, so they gain nothing there at leading order (note after Corollary II.84.5); source 24’s order- k bound is weaker than source 10’s by exactly 2^{10k+6} (note in Section F.29); source 24’s scalar constant C_3 is source 03’s κ_*^2 , and its C_2 is source 03’s $2h - h^2$ (note after Theorem II.75.1); source 24’s extraction for $\mathbb{F}_p^d$ beats source 10’s torsion budget there whenever $Q_k \leq p^{d-1}$ and $k \geq 2$ (note after Theorem II.79.5); the sharper bounds of sources 22, 25 and 27 lie below the formal bound of `lemma_15_4` in their ranges (notes after Theorems II.84.1, F.30.1 and F.32.1); source 25’s power-form exponent satisfies $K_n/E_k \rightarrow \log 2$ against the formal selection exponent (note after Theorem II.92.1); source 23’s sampling lemma measured against source 19’s theorem (note after Lemma II.103.8); the identities between sources 21 and 23 in the exponent-five theorem (notes in Section III.49). Also checked: every Lean file and line cited by sources 21–27 and by this write, at HEAD `2dfe46f44`.

Still pending. Sources 28–33 (placed by `c5513046e`), 34–39 (placed by `bf2fe146d`) and 40–46 (placed by `2dfe46f44`; both batches while this write was in progress) are in this directory and not used in this text. Their placements record, among other findings, source 29 as bounding the odd-order constants c_d^{odd} and adding a restriction theorem beside source 04’s and source 17’s, source 32 (Report276) as proving Corollary 5.8 at degree two without a scale threshold, source 33 as continuing the arrangement material of sources 14 and 15, source 34 as bounding c_4^{odd} further, source 35 as answering source 23’s question on polynomial phases, and source 38 as reducing the exponent-five problem to $\mathbb{Z}/5\mathbb{Z}$ at small u/δ , which would make source 21’s coefficient 20/9 hold in every dimension, source 41 as giving a written proof of the whole of Corollary 5.8 in every degree with no scale premise (the open `corollary_5_8`), source 42 as replacing source 26’s cube-root cap by a square root, source 44 as computing the second Whitney coefficient of the cross-section arrangement in every dimension, answering two of source 27’s questions, and source 45 as improving the upper bound for c_4^{odd} ; those findings are for the writes that add them (Section V.2).

[write] **6 October 2026; sources 28–33.** *Done in the sixth write.* Sources 28–33 (placed by `c5513046e`) are now part of the report (rows above); the fifth write’s “Still pending” paragraph describes that write only, and the title page now says thirty-three. Every delivered section of sources 28–33 is printed whole in exactly one place: their abstracts and Sections 1 in the front matter (Sections F.33–F.38, so the front matter runs to F.38), the other sections appended at the end of Parts I, II, III and V after everything printed there. In Part I, source 32’s Sections 2–5 follow source 26’s Appendix A at the end of the chapter “Corollary 5.8 without a scale threshold” (Section I.40), as that chapter’s guide planned (dated note there); no new guide. Six chapter guides were written. In Part II: the Proposition 6.1 endpoint (Section II.104; source 31’s Sections 2–5 and 9, then source 29’s Sections 5–8 as a second route), which continues source 02’s Section II.4 and source 11’s chapter; weights in the Section 7 restriction (Section II.114; source 31’s Sections 6–8 and Appendix A), which logically

follows source 10’s Sections II.20–II.21 and source 24’s chapter; restriction over a fixed field (Section II.119; source 29’s Section 9), which follows source 17’s Section II.63; Section 10 at a fixed radius and with logarithmic selection (Section II.121; source 28’s Sections 2–7 and source 30’s Sections 2–9), which continues source 12’s chapter (Section II.33) and source 07’s; and high-agreement repair of respected arrangements (Section II.136; source 33’s Sections 2–8), which continues source 14’s Section II.44 and source 15’s chapter. In Part III: the odd-order constants (Section III.66; source 29’s Sections 2–4, then source 28’s Section 8), which logically follows source 04’s Section III.3. All stand at the end of their Parts because no number printed earlier may move. Nothing printed for sources 01–27 was renumbered or relabelled; the concordance has gained the rows 28–33 (a fourth tabular).

Edits of the sixth write in delivered text of sources 28–33 (in addition to the kinds listed above): eight macros renamed where the report’s name has another meaning or prints differently (28’s `\T` as `\Tcirc`, the `\pin` of 28, 30 and 33, 29’s `\code`, 30’s `\defect`, 31’s `\Energy` and `\ip`, 33’s `\V`; they print as delivered), and unused macros not defined; source 31’s `\mathrsfs` and its column type L added to the preamble; the ‘question’ environments of sources 28, 31 and 33 (“Research question”) set as this report’s ‘Research question’, those of sources 29 and 30 (“Question”) as its ‘Question’, each numbered within its section (28, 30 and 33 numbered them on their own counter; their indices are kept); source 28’s `\cleveref` references written out as `\cleveref` prints them with its options (lowercase names; “equation (4.3)”, “lemmas 3.2, 3.3 and 4.1”); the three figure paths renamed to the shipped files; the title-page statements of sources 28, 29, 30 and 33 printed in their front-matter sections. Sources 30 and 32, like sources 02, 03, 06, 11, 12, 13 and 26, number their equations consecutively through the manuscript; here they are numbered within sections, and every other statement and equation keeps its index within its section. Tables and figures are numbered within their sections, as for sources 06–27. No write remark was needed; the counter would continue at W5.

Proved or checked at the sixth write, each inside a dated note with its argument: the intake’s certificate on $\mathbb{Z}/7\mathbb{Z}$, recomputed in exact integer arithmetic, gives $c_4^{\text{odd}} \geq 0.5587725$, above source 28’s 0.5587294 (note after Corollary III.70.8); source 29’s 35/8 is the best coefficient among sources 01–33 and sources 34, 35 and 45 record better ones (note after Corollary III.69.1); source 31’s sharp theorems supersede source 29’s endpoint constants under the same normalization (note after Theorem II.110.1); source 10’s progression is never shorter than source 31’s, so source 31’s restriction is source 10’s with weights, and source 10’s order-eight constant $2^{-209}\gamma^{103}$ beats source 31’s $2^{-363}c^{171}$, whose Reiher–Schoen formula is source 10’s (notes after Theorem II.116.1 and in Section II.116); source 30’s extension budget is implied by source 12’s sharp theorem, and composing source 30’s escape inequality with that theorem enlarges source 30’s fixed extension radius by a factor tending to 1.875 (notes after Theorem II.131.3 and Corollary II.132.2); source 32 implies the degree-two instance of the open `corollary_5_8` at every scale, and with source 26 the instances $k \leq 2$ (note after Theorem F.37.1); source 33’s linear branch exceeds θ^7 exactly for $\theta > 0.8045539\dots$ (note after Corollary II.137.6). Also checked: every Lean file and line cited by sources 28–33 and by this write, at HEAD 5d4718772.

Still pending. Sources 34–39 (placed by `bf2fe146d`) and 40–46 (placed by `2dfe46f44`) are in this directory and not used in this text. Their placements record, among other findings, source 34’s bound $c_4^{\text{odd}} \leq 2/3$ and source 45’s $c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$ (source 35 records the coefficient 5), source 36 as sharpening source 31’s noncommuting-subgroup gap, source 38 as giving the sharp odd-order U^3 stability modulus that source 31 leaves open, source 41 as a written proof of the whole of Corollary 5.8 in every degree that extends source 32, and source 42 as a square-root version of the variance step of sources 26 and 32; those findings are for the writes that add them (Section V.2).

[write] 6 October 2026; sources 34–39. *Done in the seventh write.* Sources 34–39 (placed by `bf2fe146d`) are now part of the report (rows above); the sixth write’s “Still pending” paragraph describes that write only, and the title page now says thirty-nine. Every delivered section of sources 34–39 is printed whole in exactly one place: their abstracts and Sections 1 in the front matter (Sections F.39–F.44, so the front matter runs to F.44), the other sections appended at the end of Parts I–V after everything printed there. In Part III, source 34’s Section 10 and source 35’s Section 6 follow source 28’s Section 8 at the end of the chapter on the odd-order constants (Section III.66), which was the last chapter of Part III, so that chapter stays together (dated note at its guide); no new guide. Ten chapter guides were written. In Part I: missing intervals at Lemmas 5.2 and 5.5 (Section I.49; source 37’s Sections 2–9 and Appendices A–B), which belongs before

the partition material of Section 5; polynomial divisibility and prescribed lengths (Section I.60; source 35's Section 7), which follows source 23's chapter and source 19's obstruction. In Part II: degree-weighted selection at Lemma 15.5 (Section II.144; source 35's Sections 2–5), which follows the Section 15 material of sources 17, 25 and 27; sharp odd-order stability (Section II.149; source 38's Section 6), which follows source 31's chapter on the Proposition 6.1 endpoint (Section II.104); shared anchors at Lemma 16.10 (Section II.151; source 39's Sections 2–11), which follows source 19's sampling sections; the sharp nonuniform transport threshold (Section II.162; source 34's Appendix A), which follows source 07's chapter (Section II.11). In Part III: prime torsion (Section III.73; source 34's Sections 2–9), which follows source 09's support-lattice material; quotient reduction in exponent five (Section III.82; source 38's Sections 2–5), which follows source 21's chapter (Section III.49); six-term defects (Section III.87; source 38's Section 7), which follows source 22's chapter (Section III.60). In Part IV: symmetry certification (Section IV.15; source 36's Sections 2–10), beside source 16's chapter on Proposition 17.7. All stand at the end of their Parts because no number printed earlier may move. Nothing printed for sources 01–33 was renumbered or relabelled; the concordance has gained the rows 34–39 (a fifth tabular).

Edits of the seventh write in delivered text of sources 34–39 (in addition to the kinds listed above): fifteen macros renamed or reassigned where the report's name has another meaning or prints differently (34's `\file` as `\filethirtyfour`, 39's `\file` as the report's `\filename`, the `\pin` of 35, 36, 37 and 39 and the `\shortpin` of 35 and 36, 36's `\Q` and `\R`, 37's `\T`, `\Prob` and `\code`, 38's `\Q` as `\Qthirtyeight` and its `\ip` as `\ipthirtyone`; they print as delivered), and unused macros not defined; the package `bm` and the column types `Y` and `P` added to the preamble, and room for four-digit page numbers in the contents; the 'question' environments of sources 35, 36 and 39 ("Research question") set as this report's 'Research question', that of source 37 ("Question") as its 'Question', each numbered within its section (36 and 37 numbered them on their own counter; their indices are kept); the delivered `\part` headings of sources 35 and 38 printed, unnumbered, at the head of the sections they introduce (a `\part` would have created report Parts VI and beyond); source 38's figure path renamed to the shipped file; the title-page statements of sources 35, 36, 37 and 39 printed in their front-matter sections, and source 34's abstract, which stands on its title page, printed as its abstract. No source uses `cleveref`. Every statement and equation keeps its index within its section. Tables and figures are numbered within their sections, as for sources 06–33. No write remark was needed; the counter would continue at W5.

Proved or checked at the seventh write, each inside a dated note with its argument: source 37's Corollary I.55.2 implies the conclusion of the formal `lemma_5_2`, and its Corollary I.57.3 the bound of `lemma_5_5` for $t \geq \text{weylThreshold } k$, given the formal Weyl input (notes there); source 34's $c_4^{\text{odd}} \leq 2/3$ is the best upper bound among sources 01–39, against $5^{-1/4}$ (35), $(8/35)^{1/4}$ (29) and $\sqrt{160/323}$ (28), with source 45's $(832/4795)^{1/4}$ placed for a later write (note after Theorem III.71.1); source 34's Appendix A threshold is source 07's sufficient threshold, so its sharpness answers source 07's question (note after Theorem II.163.1); source 38's odd-order modulus is compatible with source 31's all-group profile $\varepsilon/4 + 5\varepsilon^2/16 + \dots$ (note after Theorem II.150.11); source 35's frontier gives $1/(7\nu + 2)$ for quadratics and source 23's $1/(3\nu + 2)$ in degree one, for prescribed lengths only (note after Corollary I.61.6); source 39's closed profile formula holds in all 2,835 small cases for its rule and fails in 199 for source 19's rule, with the exact counterexample $m = 2$, $r = 1$, $t = 2$, $q = 2$ (note after Theorem II.155.3). Also checked: every Lean file and line cited by sources 34–39 and by this write, at HEAD 89c1cbc38, and the provenance records of the six sources (56 blob and hash records, all matching).

Still pending. Sources 40–46 (placed by 2dfe46f44) are in this directory and not used in this text. Their placement records, among other findings, source 40 as the one-stage Fejér selection on which source 35 builds (a forward reference from source 35), source 41 as a written proof of the whole of Corollary 5.8 in every degree, source 42 as a square-root version of the variance step of sources 26 and 32, source 44 as giving the second Whitney coefficient that sources 25 and 27 ask for, source 45 as recording $c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$, and sources 43 and 46 as characteristic-two counterparts of source 36; those findings are for the writes that add them (Section V.2).

[write] 7 October 2026; sources 40–46. *Done in the eighth write.* Sources 40–46 (placed by 2dfe46f44) are now part of the report (rows above; source 46 records two blob identifiers and no commit, source 40 neither); the seventh write's "Still pending" paragraph describes that write only, and the title page now says

forty-six and gives the dates 6–7 October 2026 (sources 42 and 45 are dated 7 October 2026 on their title pages, and this write was completed on that day). Every delivered section of sources 40–46 is printed whole in exactly one place: their abstracts and Sections 1 in the front matter (Sections F.45–F.51, so the front matter runs to F.51), the other sections appended at the end of Parts I–V after everything printed there. Six chapter guides were written. In Part I: Corollary 5.8 in every degree and the quadratic variance scale (Section I.62; source 41’s Sections 2–7, then source 42’s Sections 2–7), which belongs at the end of the chapter of sources 26 and 32 (Section I.40). In Part II: one-stage Fejér selection (Section II.164; source 40’s Sections 2–9), which belongs before source 35’s degree-weighted selection (Section II.144), the chapter that builds on it; the grid core of Lemma 15.4 (Section II.173; source 44’s Sections 2–10), after the degeneracy material of sources 18, 22, 25 and 27 (Section II.81); energy levels and the coset range (Section II.183; source 45’s Sections 2–3), after source 14’s near-maximal energy material. In Part III: the U^4 gap (Section III.89; source 45’s Sections 4–5), which continues the chapter on the odd-order constants (Section III.66). In Part IV: exact binary cubic norms (Section IV.25; source 46’s Sections 2–9, then source 43’s Sections 2–9 as a second route), beside source 36’s chapter (Section IV.15). All stand at the end of their Parts because no number printed earlier may move. Nothing printed for sources 01–39 was renumbered or relabelled; the concordance has gained the rows 40–46 (a sixth tabular).

Edits of the eighth write in delivered text of sources 40–46 (in addition to the kinds listed above): four macros renamed where the report’s name has another meaning or prints differently (43’s `\shortpin` as `\shortpinfortythree`, 44’s `\T`, `\card` and `\sha` as `\Tfortyfour`, `\cardtwentyfive` and `\shafortyfour`; they print as delivered), and unused macros not defined; the ‘question’ environments of sources 40, 43 and 46 (“Research question”) set as this report’s ‘Research question’, those of sources 44 and 45 (“Question”) as its ‘Question’, each numbered within its section (43 numbered its questions 1–10 on their own counter, 45 within sections on their own counter, whose indices are kept); source 46’s remarks set in this report’s remark style; source 45’s nine `\input` files printed in place and its BibTeX citations (`references.bib`, style `alpha`) pointed into this report’s bibliography; the closing group-opening lines before the bibliographies of sources 43 and 44 dropped with them; the title-page statements of sources 43, 44, 45 and 46 printed in their front-matter sections. No source uses `\cleveref` or a `\part`. Every statement and equation keeps its index within its section. Tables and figures are numbered within their sections, as for sources 06–39. *Source 41’s Sections 2 and 5 are source 32’s Sections 2 and 4, character for character* (a line diff finds no difference); they are printed again, with a note, because source 41’s later sections cite their labels and no proof is replaced by a pointer in this report. No write remark was needed; the counter would continue at W5.

Proved or checked at the eighth write, each inside a dated note with its argument: source 41’s Theorem F.46.1 implies the formal `corollary_5_8` in every degree ($D_k \leq K_k$, $\alpha/3 \geq \alpha/16$), and its Theorem I.65.3 with Lemma I.65.4 (Lemma I.65.1 at degree one) supplies the hypothesis `ThresholdFreePolynomialLocalization` of the repository’s conditional theorem, on paper (notes there); the formal statement itself was closed in the repository during this write (`b548d998e`) by its own finite localization, and is now exact; source 42’s affine theorem replaces only the modulus cap $(N/8)^{1/3}$ of source 26’s Theorem I.41.1 by $\sqrt{N}/2$, never smaller, the cube-root term in r being unchanged (note after Theorem I.74.1); source 45’s coefficient $4795/832$ is the best among sources 01–46, so $0.5587725 \leq c_4^{\text{odd}} \leq 0.6454070$, and its five-point lower example and the punctured-coset equality case $\mathbb{Z}/36\mathbb{Z} \setminus 9\mathbb{Z}/36\mathbb{Z}$ were recomputed exactly (notes after Theorem III.90.1 and Theorem II.185.4); the four claims of source 43 that the intake left unchecked were checked (note in Section F.48). Also checked: every Lean file and line cited by sources 40–46 and by this write, at HEAD `6d4942222`, the provenance records of the seven sources (30 blob and hash records, all matching), and Tidor’s displayed exponents (his Lemma 4.1 with δ^8 and Proposition 4.6 with codimension $8 \log_2(1/\delta)$, in the arXiv HTML of version 2).

Still pending. Sources 47–59 were placed in this directory while this write was in progress, by `f8bc5e2ec` (“Place batch 124”, sources 47–53) and `7dc1ee29d` (“Place batch 125”, sources 54–59); they are not used in this text. Their placements record, among other findings, source 47 as answering source 43’s quartic question (the maximum $1 - (d+1)/2^d$ for the canonical tensor C_d in every degree and dimension, so $11/16$ for every n), source 49 as recording $\|f\|_{U^4}^{16} \geq (29/4) \|f\|_{U^2}^{16}$, that is $c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$, which supersedes source 45’s coefficient, source 57 as completing the third torsion-free energy level and the bounded Schur deficits asked by source 45, and as a written proof of `theorem_13_12` through source 40’s Fejér selection, source 48 as a constructive form of the Fejér selections of sources 17 and 40, and source 50 as a written

counterexample to the catalogue statement `lemma_16_10`; those findings are for the writes that add them (Section V.2).

[write] 7 October 2026; sources 47–53. *Done in the ninth write.* Sources 47–53 (placed by `f8bc5e2ec`) are now part of the report (rows above); the eighth write’s “Still pending” paragraph describes that write only, and the title page now says fifty-three. That paragraph’s description of source 48, from its placement, as a constructive form of the Fejér selections of sources 17 and 40 overstated its scope: source 48 treats source 17’s restriction at Lemma 9.3 only (note in Section F.53). Every delivered section of sources 47–53 is printed whole in exactly one place: their abstracts and Sections 1 in the front matter (Sections F.52–F.58, so the front matter runs to F.58), the other sections appended at the end of Parts II–V after everything printed there. Six chapter guides were written. In Part II: collision-aware derandomization (Section II.186; source 48’s Sections 2–9), which belongs after source 17’s chapter on Section 9 (Section II.57); a first chapter on Section 13 (Section II.195; source 53’s Sections 2–8, then source 51’s Sections 2–10 and source 52’s Sections 2–9); two-cube degeneracy (Section II.220; source 49’s Sections 2–6), which continues source 22’s two-cube theorem in the Lemma 15.4 chapter (Section II.81); Lemma 16.10 (Section II.226; source 50’s Sections 2–9, then source 49’s Section 8), after source 39’s chapter (Section II.151). In Part III: two-harmonic dual tests (Section III.92; source 49’s Section 7), which continues the odd-order chapter (Section III.66) and source 45’s chapter (Section III.89). In Part IV: Boolean obstruction energies (Section IV.42; source 47’s Sections 2–9), which continues source 43’s part of the binary chapter (Section IV.25). All stand at the end of their Parts because no number printed earlier may move. Nothing printed for sources 01–46 was renumbered or relabelled; the concordance has gained the rows 47–53 (a seventh tabular).

Edits of the ninth write in delivered text of sources 47–53 (in addition to the kinds listed above): macros renamed or reassigned where the report’s name has another meaning or prints differently (47’s `\F`, `\Torus`, `\Sym` and `\ip` as the report’s `\Ftwo`, `\Tcirc`, `\sym` and source 31’s `\ipthirtyone`, the same definitions; 48’s `\cost` and `\pin` as `\costfortyeight` and `\pinfortyeight`; 50’s `\code` as the report’s `\filename`, the same definition, and its `\pin` as `\pinfifty`; the `\code` of 51 and 52 as `\codefiftyone`, 53’s as `\codefiftythree`; 51’s one-argument `\len` and its `\st` as `\lenfiftyone` and `\stfiftyone`, 53’s `\st` as `\stfiftythree`; the `\pin` of 51 and 53 as `\pinfiftyone` and `\pinfiftythree`; they print as delivered), and unused macros not defined; the ‘question’ environments of sources 47, 48 and 50 (“Research question”) set as this report’s ‘Research question’, that of source 49 (“Question”) as its ‘Question’, each numbered within its section (47 and 50 numbered them on their own counter; their indices are kept); source 49’s `cleverref` references written out as `cleverref` prints them with its options (lowercase for `\cref`, capitalized for `\Cref`; “equation (x)” would be its form for equations, which it does not reference this way), and the package `pgfplots` (compat level 1.18) added for its plot; source 51’s appendix group, opened before its appendix heading together with a `samepage` environment, opened after the heading instead, `samepage` dropped; source 50’s listings set in this report’s listing style; the title-page statements of sources 47 (scope and status), 48 (keywords), 49 and 50 (status) printed in their front-matter sections. Sources 51, 52 and 53, like sources 26, 30, 32 and others before them, number their equations consecutively through the manuscript; here they are numbered within sections, and every other statement keeps its index within its section. Tables and figures are numbered within their sections, as for sources 06–46. No source uses a `\part`. No write remark was needed; the counter would continue at W5.

Proved or checked at the ninth write, each inside a dated note with its argument: source 49’s coefficient $29/4$ against the earlier coefficients, so $0.5587725 \leq c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ among sources 01–53 and $c_5^{\text{odd}} \leq 29/50$ below source 34’s β_5 (note after Theorem III.93.1); the exact expansion of source 22’s L_k against source 49’s second coefficient and exact formulas, and $M_3 = 6258$ (note after Theorem II.222.4); source 47’s canonical energies $\mathcal{E}_{C_d}(1) = 1 - (d+1)/2^d$ for $d \leq 10$ and its threshold table (note in Section F.52); the constants of source 50’s counterexample; source 51’s repair exponent 2γ against source 53’s 4γ (note after Corollary II.211.4); the exponent arithmetic of source 52 and the scalar identities and elementary inequalities of source 53’s threshold collapse. Also checked: every Lean file and line cited by this write at HEAD `06b30ca52`, the provenance records of the seven sources (79 hash, blob and pin records, all matching), the Lean files they cite that changed after their pins, and the repository’s formal work on their subjects since then (Section V.1).

Still pending. Sources 54–59 were placed in this directory by `7dc1ee29d` (“Place batch 125”) while the

eighth write was in progress; they are not used in this text. Their placement records, among other findings, source 54 as a single global quadratic phase at Section 8, source 55 as endpoint stability in every U^d (agreeing at $d = 3$ with sources 31 and 38), source 56 as a relative affine repair beside source 11’s chapter, source 57 as a written proof of `theorem_13_12` through source 40’s Fejér selection (the repository has since closed that statement formally, `d48e8af00`, crediting inputs to source 57’s route in `FORMALIZATION_STATUS.txt`) and as the complete third torsion-free energy level asked by source 45, source 58 as Lemma 9.3 on every cyclic group, and source 59 as character-affine refinements of Section 7 and source 11’s chapter. Sources 60–64 and 65–69 were placed by `39d3295c4` and `06b30ca52` (“Place batch 126” and “Place batch 127”) while this write was in progress; they are not used in this text either. Their placement records source 60 as an all-order bound for the vertical difference failure answering source 33’s Questions 2–3, source 61 as an exact, threshold-free restriction on arbitrary finite abelian groups, source 62 as restriction on every cyclic modulus by unit-dilation averaging (beside source 58), source 63 as $\lim_d c_d^{\text{odd}} = 1/3$, answering the limit questions of sources 28 and 29, source 64 as binomial separation for Gowers’s Section 4 conjectures, sources 65 and 69 (independently) as the extremal energy $1 - (d + 1)/2^d$ of nonintegrable symmetric d -linear tensors on $\mathbb{F}_2^n$ in every degree $d \geq 3$, answering source 47’s Questions 1–2, source 66 as endpoint partitions and the fourth energy level beside source 57, source 67 as exact local profiles at source 11’s affine chapter, and source 68 as sharp energy and restriction bounds (coset existence up to $\varepsilon = 1/4$; Lemma 9.3 on every cyclic group with exponent 53). Those findings are for the writes that add them and were not checked here (Section V.2).

[write] **7 October 2026; sources 54–59.** *Done in the tenth write.* Sources 54–59 (placed by `7dc1ee29d`) are now part of the report (rows above); the “Still pending” paragraphs of the eighth and ninth writes describe those writes only, and the title page now says fifty-nine. Every delivered section of sources 54–59 is printed whole in exactly one place: their abstracts and Sections 1 in the front matter (Sections F.59–F.64, so the front matter runs to F.64), the other sections appended at the end of Parts II and V after everything printed there. Seven chapter guides were written, all in Part II, in the order of Gowers’s sections: the polynomial-phase endpoint (Section II.236; source 55’s Sections 2–11, then source 57’s Section 5), which continues the chapters of sources 31 and 29 (Section II.104) and 38 (Section II.149); affine rigidity at the Section 6 endpoint (Section II.248; source 59’s Sections 7–10, then source 56’s Sections 2–9), which continues source 11’s chapter (Section II.22); character separation at Section 7 (Section II.261; source 59’s Sections 2–6), after the chapters of sources 24 and 31 (Sections II.72 and II.114); one quadratic phase at Proposition 8.1 (Section II.267; source 54’s Sections 2–9), the report’s first chapter on Section 8; restriction on every cyclic group at Lemma 9.3 (Section II.276; source 58’s Sections 2–11), after the chapters of sources 17, 29 and 48 (Sections II.57, II.119 and II.186); Theorem 13.12 (Section II.287; source 57’s Sections 2–4), which continues the ninth write’s Section 13 chapter (Section II.195); and the third energy level (Section II.291; source 57’s Section 6), which continues source 45’s chapter (Section II.183). All stand at the end of Part II because no number printed earlier may move. Nothing printed for sources 01–53 was renumbered or relabelled; the concordance has gained the rows 54–59 (an eighth tabular).

Edits of the tenth write in delivered text of sources 54–59 (in addition to the kinds listed above): macros renamed or reassigned where the report’s name has another meaning or prints differently (54’s `\ee`, the letter e, as the report’s `\e`; 55’s `\U` as `\Ufiftyfive`; 56’s `\Q` as source 49’s `\Qcal`; 58’s `\GG` and `\commit` as `\GGfiftyeight` and `\commitfiftyeight`; the `\repo` of 58 and 59, in small capitals, as `\repsc`; they print as delivered), and unused macros not defined; source 57’s ‘question’ environment (“Research question”, numbered 1–12 on its own counter) set as this report’s ‘Research question’, numbered within its section with its indices kept; the ‘question’ environments of sources 54, 55, 56, 58 and 59 (“Question”) set as this report’s ‘Question’ (58’s and 59’s numbered within their sections on their own counter; their indices are kept); source 57’s ‘Input theorem’ environment and source 58’s unnumbered ‘Theorem A’ and ‘Theorem B’ added; source 56’s thirteen section files, `\input` by its main file, printed in that order; the `cleveref` references of sources 55 (options `nameinlink`, `noabbrev`) and 57 (`capitalise`, `noabbrev`) written out, 55’s as `cleveref` prints them (lowercase), 57’s with the capitalized name of the statement referred to, where its delivered PDF prints “Theorem” for every statement on its shared counter (for example “Theorem 2.2” for Lemma 2.2); the formatting lines that open the delivered bibliographies of sources 55 and 59 (a small-type group) dropped with them; the title-page statements of sources 55, 56 and 59 (status), 57 (repository baseline) and 58 (keywords) printed in their front-matter sections; source 58’s figure read from `figures/58-torsion-restriction-kernel_exponents.pdf`. All six number their equations within sections, and every statement and equation keeps its index within its

section. Tables and figures are numbered within their sections, as for sources 06–53. No source uses a `\part`. No write remark was needed; the counter would continue at W5.

Proved or checked at the tenth write, each inside a dated note with its argument: the interface of source 57’s Theorem F.62.1 against the catalogue’s `theorem_13_12`, every exponent of its chain in exact integers, and the formal proof’s route and constants, which reach density $(\alpha/2)^{2^{42}}$ where source 57 reaches $(\alpha/2)^{2^{41}}$ (note after Theorem II.290.3); a closed form of source 57’s threshold $N_0(\alpha)$ from source 53’s Theorem F.58.1, which answers its Research question V.216.8 (note after Remark II.290.4); source 55’s coefficient identities and table, and its $d = 3$ coefficients against sources 31 and 38 (note after Theorem II.237.1); source 54’s constant $27/4$ and the comparison with Proposition 8.1 and the formal statement (note after Theorem F.59.1); the rational constant chain of source 56; source 58’s five integer certificates and its comparison with Lemma 9.3; source 59’s cutoff arithmetic ($\Psi(1/16) = 435/2048 < 2/9$, $\min_{[1/7, 1/3]} \Psi = 132/343$, $1296 > 1266$) and the order-five probability $3/5$ in its remark on Green and Tao’s Lemma 8.3. Also checked: every Lean file and line cited by this write at HEAD `fe7165a3c`, the provenance records of the six sources, and the repository’s formal work on their subjects since their pins (Section V.1).

Still pending. Sources 60–64 and 65–69 were placed by `39d3295c4` and `06b30ca52` (“Place batch 126” and “Place batch 127”) before this write began, and sources 70–75, 76–81 and 82–88 by `b788eda88`, `487859db5` and `bbde2b873` (“Place batch 128” to “Place batch 130”) while it was in progress; none of them is used in this text. Among the findings their placements record, those that bear on sources 54–59 are: source 61 finds source 59’s restriction exponent $(2m + 1) \log_2(5m)$ independently and gives an exact, threshold-free restriction on arbitrary finite abelian groups; sources 62 and 68 give Lemma 9.3 on every cyclic modulus by other methods, 68 with exponent 53 (source 58: 54); source 66 gives the endpoint partitions and the fourth energy level beside source 57’s Section 6, uncited (source 57 has the better third-level cutoff, $m \geq 7$); source 67 gives the non-affine energy maximum from $N \geq 13$, with 13 optimal for elementary abelian 3-group targets (source 59’s $\mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/3\mathbb{Z}$ example, uncited there); source 70’s headline is the printed Theorem 13.12, already proved in writing by source 57 and formally, with a near-maximal branch on every cyclic modulus; source 72 gives a near-maximal frequency map on every finite abelian group beside source 57’s Section 5; source 73 gives an explicit entry threshold $\varepsilon \leq 10^{-d}$ and radius $q^{-5/2}$ for source 55’s sharp expansion, which would answer source 55’s Question V.208.4; and sources 76–83 give sharp hereditary energy-rigidity constants for full-domain maps, to be printed after source 56 in the affine chapter. Those findings are for the writes that add them and were not checked here (Section V.2).

[write] 7 October 2026; sources 60–69. *Done in the eleventh write*. Sources 60–69 (placed by `39d3295c4` and `06b30ca52`) are now part of the report (rows above); the “Still pending” paragraphs of the earlier writes describe those writes only, and the title page now says sixty-nine. Every delivered section of sources 60–69 is printed whole in exactly one place: their abstracts and Sections 1 in the front matter (Sections F.65–F.74, so the front matter runs to F.74), the other sections appended at the end of Parts I–V after everything printed there. Eleven chapter guides were written, in the order of Gowers’s sections within each Part: unequal Fourier budgets at Lemmas 5.2 and 5.5 (Part I, Section I.75; source 67’s Section 2), after source 37’s chapter; in Part II, affine rigidity in characteristic three (Section II.293; source 67’s Sections 3–4), the energy gap, the coset range and the fourth energy level (Section II.296; source 68’s Sections 2–4, then source 66’s Sections 2–11), Lemma 9.3 on every modulus (Section II.310; source 68’s Sections 5–7, source 62’s Sections 2–9, source 61’s Sections 2–9 and source 60’s Section 9), the arrangement polynomial of every order (Section II.331; source 60’s Sections 3–8) and the interpolation count at Lemma 16.10 (Section II.338; source 67’s Section 5); in Part III, the odd-order limit (Section III.94; source 63’s Sections 2–5, then source 60’s Section 2), Gowers’s Section 4 conjectures (Section III.100; source 64’s Sections 2–8) and norm layers without a degree cap (Section III.108; source 67’s Section 6); in Part IV, Boolean phase integration in every degree (Section IV.51; source 65’s Sections 2–8, then source 69’s Sections 2–9) and the face cover at Proposition 17.7 (Section IV.67; source 63’s Section 6). All stand at the ends of their Parts because no number printed earlier may move; each guide says which earlier chapter it continues. Nothing printed for sources 01–59 was renumbered or relabelled; the concordance has gained the rows 60–69 (a ninth tabular).

Edits of the eleventh write in delivered text of sources 60–69 (in addition to the kinds listed above): macros renamed or reassigned where the report’s name has another meaning or prints differently (60’s `\V` and `\pin` as `\Vcal` and `\pinsixty`; the `\repo` of 61 as `\repsc`, as for 58 and 59; 62’s `\ee`, `\U` and `\R` as `\eexp`, `\Urm`

and `\Rcal`; the circle $\mathbb{R}/\mathbb{Z}$, written `\T`, `\TT` and `\Torus` by 63, 65 and 69, as `\Tquot`; 64’s one-argument `\repo` as `\repopath`; the $\mathbb{F}$ of 65 and 69 ($\mathbb{F}_2$) as the report’s `\Ftwo`; 65’s `\pin` as `\pinsixtyfive`; 66’s `\E`, `\D`, `\ep`, `\card` and `\wt` as `\Erm`, `\Drum`, the report’s `\e`, `\cardbars` and `\wtop`; 68’s `\sym` and `\cost` as the report’s `\syndiff` and as `\costA`; they print as delivered), and unused macros not defined; the ‘question’ environments of 63 (“Research question”, on the shared counter), 65 and 69 (“Research question”, on their own counters) and 60’s ‘problem’ (“Research question”, on its own counter) set as this report’s ‘Research question’, numbered within their sections, which keeps the indices; the ‘question’ environments of 61, 62, 64, 66, 67 and 68 set as this report’s ‘Question’; source 66’s ‘conjecture’ environment added; source 64’s `cleveref` references (options `capitalize`, `nameinlink`, `noabbrev`) written out with the capitalized name of the statement referred to, where its delivered PDF prints “Theorem” for every statement on its shared counter, as source 57’s did; source 67’s and source 68’s single-file assemblies printed, with the assemblies’ comments kept inside sections and those at the joins between section files dropped, as also in the delivered single files of sources 63 and 64, which carry such join comments; the `\appendix` and `\clearpage` lines between sections dropped, as before; the post-abstract statements of every source (scope, status, keywords) printed in their front-matter sections; the figure paths renamed to the shipped files. All ten number their equations within sections, and every statement and equation keeps its index within its section; tables and figures are numbered within their sections, as for sources 06–59. No source uses a `\part`. No write remark was needed; the counter would continue at W5. One layout parameter changed: Part III now has a Section III.100, so the width of section numbers in the table of contents was raised from 2.9em to 3.0em; no number moved.

Proved or checked at the eleventh write, each inside a dated note with its argument: that source 63’s sharp bounded inequality settles source 60’s bracket for β_{odd} and the first half of its finite-cyclic question, $b(\mathbb{Z}/q\mathbb{Z}) = (q^3 - 4q + 3)/(3q^3)$ (note after Theorem III.96.4), and that its rate settles source 60’s effectivity question for c_d^{odd} and for the large-prime limit; the comparison of the five all-moduli restriction theorems of sources 58, 60, 61, 62 and 68, setting by setting, with the exact crossover $\alpha^{22}\eta^{23}(1 - \eta) < 2^{-139}$ between the retentions of sources 62 and 68 (note in Section II.310), and that the theorems of sources 61, 62 and 68 imply `lemma_9_3` on paper (61 also `corollary_9_4`, with $N_0 = 1$); the integer certificates of sources 61 and 68; the misprinted list of seven numerators for eight orders in source 61’s root-grid lemma (the eighth value is 3, and the lemma holds as stated; note there); that source 64’s dimension barrier is source 10’s exact threshold at degree $k - 1$; the class counts of sources 65 and 69 against each other; source 66’s thresholds against source 57’s $(3r + 1)$ against $5s - 1$, with $2r + 2 \leq N(r) \leq 3r + 1$; source 68’s thresholds against sources 14 and 45 and its certificates; the order-thirteen comparison of sources 59 and 67; source 63’s face-cover denominators and the comparison with source 16’s coefficients. Also checked: every Lean file and line cited by this write at HEAD `fb9602e55`, the provenance records of the ten sources, and the repository’s formal work on their subjects since their pins (Section V.1).

Still pending. Sources 70–75, 76–81 and 82–88 were placed by `b788eda88`, `487859db5` and `bbde2b873` (“Place batch 128” to “Place batch 130”) before this write began, and sources 89–94, 95–96, 97–98 and 99–100 by `39beb65f2`, `09f5b745a`, `9599d8566` and `53eedc74c` (“Place batch 131” to “Place batch 134”) while it was in progress; none of them is used in this text. Their placements record findings for these items: II.55 (source 70: the printed Theorem 13.12, already proved in writing by source 57 and formally, with a near-maximal branch on every cyclic modulus and a lower-range exponent conditional on Lau’s recurrence theorem; source 71: the Stage 13.7 obstruction of source 52 on fixed mixed prime support); II.2 and II.58 (source 72: a near-maximal frequency map on every finite abelian group beside source 57’s Section 5; source 73: an explicit entry $\varepsilon \leq 10^{-d}$ and radius $q^{-5/2}$ for source 55’s expansion, which would answer source 55’s Question V.208.4); II.41 and II.56 (sources 74 and 87: the packaged Lemma 16.10 premises and the contextual lift contract fail on sparse and positive-density slabs; source 75: a deletion fraction at parameter one); II.19 (sources 75–83 and 85: sharp hereditary relative-energy constants for full-domain maps, $3/4$, $19/27$, $\rho = (\sqrt{6} - 1)/(\sqrt{6} - 1 + \sqrt{\sqrt{6} - 2})$, and on $\mathbb{F}_3$ -spaces $5/8$, $53/81$, λ and $5/9$, to be printed after source 56 in the affine chapter); V.1 (sources 84 and 86: comparisons of the repository’s explicit five-term thresholds); II.32 and II.50 (source 88: product-property refinements of source 15’s chapter, with a selection corollary for source 35’s); I.20 (sources 90 and 93: the complete four-frequency missing-arc profile Φ_4 , independently, with grid corrections); II.23 and II.21 (source 93: affine recovery below graph defect $1/4$ for every abelian target, and the non-affine maximum $N^3 - 4N^2 + (10 + 2\kappa)N - (6 + 2\kappa)$ from $N \geq 13$, optimal, for every target; source 90: a sharp bound for maps $C_N \rightarrow W$, $N \geq 12$); III.1 (source 93: $c_d(F) - D(F) = O(d2^{-d/2})$)

uniformly over odd-order families, a rate that would supersede source 63's); IV.9 (sources 89, 91 and 92: the first nonclassical obstruction in odd characteristic, its repair subspaces and covers, and depth reduction); III.26 (source 94: sharp three-term counting with arbitrary homomorphisms, beside source 20's chapter); IV.10 (source 95: simultaneous integration of two Boolean cubics through the generic half-rank of their pencil); II.19, II.49 and III.1 (source 96: the affine cubic $4t - 10t^2 + 6t^3$ on the whole range $t \leq 2/3$ for every odd-order target; source 60's odd-target arrangement repair and the bracket $[15/256, 23/256]$, proved independently, in a version with no cubic term and no smallness condition; source 33's erasure-robust repair without its additive floor; $\|f\|_{U^4}^{16} \geq (843/100) \|f\|_{U^2}^{16}$, that is $c_4^{\text{odd}} \leq (100/843)^{1/4} = 0.5868716$); IV.11 (source 97: the maximizing functions of the canonical tensor for $d \geq 4$, $e(P + \theta|u|)$ with $\deg P \leq d - 1$, and dimension-free stability with optimal leading coefficient, which would answer the function questions of sources 65 and 69); II.19 (source 98: a is affine iff $r_3(a) > 5/8$, sharp, for full-domain maps between arbitrary abelian groups); II.48 (source 99: fixed-radius Bohr sets in every large prime cyclic group with exactly $2^r - 1$ runs, so $(2^r - 1)/2 \leq c_r^{\text{prime}}$, which would answer source 30's prime-order question and raise the lower rank base from $\sqrt{3}$ to 2); III.25 and III.8 (source 100: sets of density δ in every large prime cyclic group whose k -term count deviates by a nonzero constant times the squared U^d norm, so that no $o(u^2)$ counting error is possible, which would answer source 20's Questions 1 and 3). Among them, those bearing on sources 60–69: source 97 on the function questions of sources 65 and 69; source 96 on source 60's arrangement chapter (the same results, independently; source 60 arrived first and is printed in its form) and on source 67's characteristic-three range; source 93 on source 67's recovery threshold and cutoff (for every abelian target where source 67 has elementary abelian 3-groups) and on source 63's rate; sources 90 and 93 on source 67's Fourier profiles; sources 89, 91 and 92 on the odd-characteristic questions of sources 65 and 69; sources 80–83 work on $\mathbb{F}_3$ -spaces beside source 67's characteristic-three chapter (domains, where source 67 has targets); and source 74's and 87's slab obstructions concern the Lemma 16.10 interface whose affine count source 67 sharpens. Those findings are for the writes that add them and were not checked here (Section V.2).

F.2 What is proved, and where

“Proved” means proved in the cited source by an ordinary written proof, unrefereed. “External” marks a published theorem used and not reproved. Nothing here is checked by a proof assistant, except the one formalized statement marked **formal** (formalized before the sources arrived); the report gains no formal status (Section F.5). (The last clause was added after the independent check of 6 October 2026.)

Statement	Where	Sources
<i>Part I. Density transfer and phase-flat partitions</i>		
Sharp one-sided mass bound	Thm. I.9.1, Thm. I.14.1; case	02, 03; 01, 05
$\Pr(X > t) \geq ((a + t)T - at)/(a(b - t))$ for mean-zero	$a = \delta$: Thm. I.4.1,	
$X \in [-a, b]$, $T = \mathbb{E}X_+$; forced maximum $aT/(a - T)$;	Lem. I.17.3	
extremizers		
Short-arc density tail with phase parameter s ,	Thms. I.2.1, I.2.3, I.3.1	01
occupied-mass exceptions, optimal forced maximum a_* ,		
the trade-off $s_{\max}(t)$; sharp in the finite weighted model,		
not claimed for arithmetic characters		
One-cell phase envelope $C^2 \leq (1 - \rho^2)F^2 + \rho^2L^2$ and	Lem. I.1.2, I.1.1; Thm. I.8.1,	01, 02
the linear law $C \leq \rho L + (1 - \rho)S$ (real f only)	Cor. I.8.2	
Finite-knot linear-programming reduction; multiphase	Thm. I.4.2, Cor. I.4.3	01
score		
Lemma 5.15 with $ P \geq \alpha(2 + \alpha)N/(2(4 - \alpha)(M - 1))$	Cor. I.9.3, (I.14.4)	02, 03
(needs $\alpha > 0$, nonempty cells)		
Simultaneous phase-flat partition,	Thm. I.5.1, Thm. I.18.2	01, 05
$L = \lfloor (N \prod_i \varepsilon_i / 2^m)^{1/(m+1)} \rfloor$, every cell in $[L, 2L - 1]$		

Statement	Where	Sources
Fourier increment $\delta + 3\delta b\kappa/(4 - (4 - 3\delta)\kappa) \geq \delta + 3\eta/8$, $b = 1 - \delta$, length $\lfloor \sqrt{(N/2\pi) \arcsin(\eta/4\delta b)} \rfloor$; real frequency, bounded weights, any N	Thm. I.6.1	01
Fourier increment $\delta + \delta\alpha/(4\delta - \alpha)$, length $\geq \sqrt{\alpha N/(16\pi\delta(1 - \delta))}$, exact integer length, popularity, trade-off	Thm. I.17.4	05
Fourier increment with adjustable loss λ	Thm. I.15.2	03
Increment comparison: source 01's is at least 3/2 times source 05's (and source 03's at $\lambda = 1/2$) at stated parameters	[write] note after Table 3	write
Lemma 2.3 with lower constant 1/3 (optimal, also under the exact Lean hypotheses)	Thm. I.16.4	05
Corollary 2.4 with 36π and $3\alpha/4$, same upper length	Cor. I.16.6	05
Optimal simultaneous localization order $1 + (\varepsilon^m N)^{1/(m+1)}$; one-phase barrier $\ell \leq 1 + \sqrt{3r(N - 1)}$	Thm. I.18.5; Thm. I.7.2	05; 01
Spectral extraction $E \leq \sqrt{EV} + 2R(v - V)$, Gram correction, optimal width allocation, short-cell discard	Thm. I.19.2, Cor. I.19.4, Prop. I.19.3, Prop. I.20.1	05
Robust phase-to-density transfer; refinement with one global phase budget and its optimum, <i>conditional on</i> $PP(K, R_0)$ (Corollary 5.6, now formal for intervals) <i>Part I, sources 06–10 (second write)</i>	Thms. I.10.1, I.11.1, Cor. I.11.2, Prop. I.11.3	02
Exact fixed- M optimum $\min\{P_+/(b(M - 1)), S/(M - 2)\}$ for the largest high cell, with extremizers; Lemma 5.15 with both branches (strictly better than the $M - 1$ corollary for $\alpha < 1$); simultaneous $\tau(\alpha)$; signed drift and exceptional sets	Thm. I.21.1, Cors. I.21.2, I.21.3, Prop. I.21.4	06
Sharp one-sided mass bound, now in six sources	Thm. I.21.1; Prop. I.23.2; note after Thm. I.4.1	06, 10 (with 01, 02, 03, 05)
Optimal cap $M \geq \delta + \delta S/(2\delta - S)$ and variance cap $M \geq \delta + V/\delta$; correlation transfer with an exceptional set	Thm. I.23.1, Thms. I.23.3, I.23.4	10
Explicit phase partition and the density–length frontier D_γ, L_γ ; variance form of the U^2 inverse step	Lem. I.24.1, Thm. I.24.2, Cor. I.24.3, Prop. I.24.4	10
Source 01's increment is at least 3/2 times source 10's for $\gamma \leq 1/2$, at least equal for $\gamma \leq 3/(4 - \kappa)$ <i>Part I, sources 11–14 (third write)</i>	[write] note after Table 3	write
Complete simultaneous partition into lengths $v - 1$ or v , diameters $\leq \varepsilon_i$, for $1 \leq v \leq (rE_0)^{1/(q+2)}$; linear	Thm. I.26.8, Cor. I.26.9	13
Lemma 5.9 with exponent $1/(3q + 2)$ and no size threshold		
Flexible-length frontier $(q + 1)a + qb = 1$, optimal; random-slope obstruction with independent tolerances $\varepsilon_i < 1/2$ and a finite prime certificate	Thm. I.26.5, Cor. I.26.6	13 (with 01, 05)
Sharp one-sided bound, now in seven sources; joint phase–density–exception frontier (= source 01's)	Thm. I.27.1, Thm. I.28.3	13 (with 01, 02, 03, 05, 06, 10)
Fourier increment $\delta + \alpha/(4 - \alpha(2 - \delta)/(\delta(1 - \delta)))$; dominated by source 01's at the same length <i>Part I, sources 15–20 (fourth write)</i>	Thm. I.29.1; [write] note there	13; write

Statement	Where	Sources
Exact partitions whose own steps control affine frequencies: a finite certificate theorem; every balanced exponent $\gamma < 1/((q+1)(42q+1))$, conditional on Lau's preprint (threshold not numerical); weighted cell $\alpha^{32}/16$; obstruction $\gamma \leq 1/(3q+1)$	Thms. I.30.2, I.30.5, Lem. I.30.7, Cor. I.31.2	16
Simultaneous polynomial partitions into lengths $v-1$ or v with exponents polynomial in the number of phases (quadratics: $a < 1/(126q^2 + 129q + 2)$; shared leading coefficient: $1/(99q + 98)$; degree profiles $F_k a + G_k b < 1$; conditional on Maynard's theorem (numbers from Lau's preprint); source 13's Question 3 answered	Thm. I.32.2, Cors. I.32.4, I.32.5; Thm. I.34.3, Cor. I.34.5	18, 19
Volume obstruction $(qk(k+1)/2 + 1)a + kqb \leq 1$, and $(S+1)a + Db \leq 1$ for degree profiles <i>Part I, sources 21–27 (fifth write)</i>	Cor. I.33.5; Thm. I.34.7	18, 19 ($k = 1$: 13, 05)
Prescribed-length simultaneous linear partitions: the universal threshold lies between $C(S) \prod_i M_i$ and $C(S)B$ (C the semigroup conductor); $L^{q+2}/\prod_i \varepsilon_i (1 + o(1))$ for lengths $L, L+1$; the frontier $(q+2)a + qb = 1$ is optimal even with cell-dependent steps, so source 13's exponent $1/(3q+2)$ is optimal (its Question 1 answered); exact packing for the obstructions; almost-partitions	Thm. I.36.1, Cor. I.37.1, Prop. I.38.1, Thm. I.39.1, Cor. I.39.4	23 (upper bound 13)
Corollary 5.8 at degree one without a scale threshold: gain $\alpha/3$ at length $\min\{(r(1-\delta)/(128\mu))^{1/3}, (N/8)^{1/3}\}$, or $(N/M)^{1/3}/8$ for two-comparable parents; a parent-contained alternative $\sqrt{\alpha r/(24\pi\mu)}$; every cyclic modulus; implies the degree-one instance of the open formal statement ([write] note)	Thms. I.41.1, I.41.2, Cor. I.41.3	26; write
Common-step integer-box almost-cover, width $> \eta\sqrt{m}/(3D)$; exact long ordinary partitions of arbitrary proper modular parents are impossible <i>Part I, sources 28–33 (sixth write)</i>	Thm. I.44.1, Sec. I.44.1	26
Corollary 5.8 at degree two without a scale threshold: gain $\alpha/3$ at length $\max\{1, (N/M)^{1/39}/17\}$ over every cyclic group, two-comparable parents; implies the degree-two instance of the open formal statement (with source 26 the instances $k \leq 2$; [write] note); explicit recurrence $p \leq 2^{10}R^5$ with $\ \theta p^2\ < 1/R$; lossless quadratic partitions above $2^{116}L^{38}$ <i>Part I, sources 34–39 (seventh write)</i>	Thm. F.37.1, Lems. I.45.3, I.46.2, I.47.1	32 (with 26)
Exact first-window missing-arc minimax $\rho_R(a) = 1/(R - \cot(a/2) \tan((R+1)a/2))$ for $\pi < (R+1)a < 2\pi$, unique extremizer, exact mass tradeoff, stability; Lemma 5.2 with frequency $\leq \lfloor N/M \rfloor - 1$ and coefficient $ A M/(N-M)$ for any weights (implies the formal statement; [write] note); recurrence exponent $(1 - 1/(4k))/(2^{k-1} + 1)$ from the formal Weyl input (implies the formal Lemma 5.5 bound; [write] note)	Thm. I.50.1, Cors. I.55.2, I.57.3, Thm. I.56.3	37
Explicit polynomial phases forcing $D \mid h$ on every flat progression; exact thresholds $D C(S)$; necessary frontier $(W+2)a + \sum_i d_i b_i \leq 1$ for prescribed lengths (quadratics: $1/(7\nu+2)$)	Thms. I.61.2, I.61.4, Cor. I.61.6	35 (with 23, 19)

Statement	Where	Sources
<i>Part I, sources 40–46 (eighth write)</i>		
Corollary 5.8 in every degree, no scale premise: gain $\alpha/3$ at length $\max\{1, (N/M)^{1/D_k}/8\}$, $D_1 = 4$, $D_k = K_k/4 + 1$; explicit recurrence $p \leq 2^{A_k} R^{B_k}$ for every $R \geq 2$; lossless localization above $2^{c_k} L^{e_k} \leq (2L)^{K_k/4}$; compact bound $2^{-(k+8)}(N/M)^{1/(12\mathcal{P}_k+1)}$; retained bounds of sources 26 and 32 and a cubic $(N/M)^{1/4096}/8$ (implies the formal statement, [write] note; that statement is now exact, by the repository’s own proof)	Thm. F.46.1, Thms. I.64.3, I.65.3, Cors. I.68.2, I.68.5	41 (with 26, 32)
Arithmetic proper-progression increment $\delta + \frac{1-\delta}{L} \frac{\max\{q-L, 0\}}{q-1}$ on every cyclic group; gain $(1-\delta)/(2L)$ for $N \geq 2(L-1)^2 + 1$; sharp $1 - 1/C$ under $N \geq CL^2$ (prime-quotient extremizers); sliding-window gain $> \delta(1-\delta)/(2L)$ for $N > L(L-1)$; affine cap $\sqrt{N}/2$ in place of $(N/8)^{1/3}$	Thm. F.47.1, Cor. F.47.2, Thms. I.72.3, I.73.1, I.74.1	42 (with 26, 41)
<i>Part I, sources 60–69 (eleventh write)</i>		
Lemma 5.2 with unequal budgets: exact missing-arc mass $-1/((1-\cos a) \mathbf{e}^T L(\boldsymbol{\varepsilon})^{-1} \mathbf{e})$ on an explicit neighbourhood of every point of source 37’s equal-budget branch; the old certificate’s optimality polytope; Hessian rank $\lfloor (R-1)/2 \rfloor$; aligned grids to order N^{-2}	Thms. I.76.2, I.76.4, I.76.5, I.76.6, I.76.8	67 (with 37)
<i>Part II. The inverse step</i>		
$\mathcal{T}_w(f, g)^4 \leq \min\{m(f)^4 \mathcal{U}(g), m(g)^4 \mathcal{U}(f)\} \mathcal{E}(w)$ on every finite abelian group, coefficient 1 optimal; $\ f\ _{U^3} \leq \ f\ _2$; graph and balanced forms; spectral certificate; deficits; alignment; rigidity	Thm. II.2.1 and Sections II.2–II.4	02
Lemma 7.4 with $ K \geq 7\delta^3 n/(3\sqrt{6})$, same threshold and 90%; optimal scalar tangent κ_* ; weighted tuples; one-witness trade-off	Thm. II.5.2, Section II.5	03
Elementary Proposition 7.3: $ A' \geq c A /8$, $ A' - A' \leq 2^{19} c^{-9} A $ (weaker than Reiher–Schoen)	Thm. II.6.3	03
Exact energy interpolation (M1)–(M7); Lemma 9.2 strengthened to $\gamma^7 \beta^{15} p^{15}$	Thm. II.8.1, Cor. II.8.2, Prop. II.8.3	04
Polynomial restriction over $\mathbb{F}_p$; order eight: $2^{-95} \gamma^{112} \eta^{15} B ^{15}$ under $p \geq 2^{102} \gamma^{-112} \eta^{-16} \beta^{-1}$; method limit (prime fields only)	Thm. II.9.1, Cor. II.9.5, Prop. II.9.6	04
Arrangement interpolation with profile σ ; Lemmas 14.8, 12.4 with $\beta^{7.4k+1-6}$, β^{106}	Thm. II.10.1	04
<i>Part II, sources 06–10 (second write)</i>		
Coupled-path rigidity with transport profiles; Freiman order k when $2[1 + (k-1)\kappa]\delta < 1$; sharp threshold $\delta < 1/(2k)$ for constant fibres	Thm. II.15.1, Cor. II.15.3, Thm. II.16.1	07
Lemma 10.9 with Freiman order five under exactly the Lean hypotheses; order nine with the ambient package (not among the Lean hypotheses); allowances 120^{-5} , 75^{-5}	Cors. II.17.3, II.17.6, Props. II.17.4, II.17.7	07
Root-law linear program; collision-to-mode exactification; weighted laws; nonabelian groups	Thm. II.18.3, Props. II.19.1, II.19.2, Thm. II.19.3	07
Proposition 7.3 with $(c/4, 2^{10} c^{-5})$ on any abelian group	Thm. II.20.1	10
Lemma 7.5 with $ B /(k(h_k - 1) + 1)$; Corollary 7.6 with $2^{-209} \gamma^{103}$; general finite codomains	Thm. II.21.1, Cor. II.21.2, Thm. II.21.3	10

Statement	Where	Sources
<i>Part II, sources 11–14 (third write)</i>		
Sharp cubic profiles $\varepsilon(f) \geq 4t - (10 + 2\kappa)t^2 + (6 + 2\kappa)t^3$ with coset extremizers; correction for $\varepsilon < 1/9$; inverse series; exact largest non-affine graph energy for $ G \geq 34$ and for all $\mathbb{F}_p^n \rightarrow \mathbb{F}_p^m$	Thms. F.16.1, F.16.2, F.16.3	11
The cutoff 34 cannot be dropped for general groups ($\mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/2\mathbb{Z}$: $1296 > 1288$)	Rem. W2	write
High-correlation refinement of Proposition 6.1 ($\alpha^4 > 8/9$; every N , see the note)	Cor. II.32.2	11
Sharp Freiman extension $3\varepsilon + 2\kappa < 1$ with boundary counterexamples; global extension above $2/3$; sharp collision-to-mode $\chi(\eta)$; directional repair	Thms. II.35.2, II.35.3, Cor. II.35.5, Lem. II.36.1, Thm. II.36.2	12
Radius selection $\rho/(8d)$; fourth-moment smoothing $2^{-74}\alpha^9/\pi$; Theorem 10.13 with radius $\alpha^{19}/(3\pi 2^{149})$ and $\alpha^5 X /20000$ (the α^5 is already proof-level in the formal proof)	Thm. II.37.2, Thm. II.38.1, Thm. II.40.2, Cor. II.40.3	12
$E(A) \geq (1 - \varepsilon)m^3$, $\varepsilon \leq 1/100$, forces a unique coset with $ A \triangle C \leq \tau(\varepsilon)m$, sharp; product cosets for cube collisions; equality with labels; arrangement-moment equality	Thms. II.43.2, II.43.3, II.43.5, II.44.1, II.44.4	14
Near-extremizers of source 11 lie near a coset at the scale η/t^3 (source 11's Question 3 in substance)	Rem. W3	write, from 11, 14
<i>Part II, sources 15–20 (fourth write)</i>		
Collision-sensitive amplification $R_2^{d-1} \leq C^{d-2}R_d$, cubic cap, density-sensitive Lemma 14.7 (as source 04's (A5)–(A11) and source 14's Theorem 5.1)	Thm. II.47.3, Lem. II.48.1, Cor. II.48.2	15 (with 04, 14)
Collision-parametrized product step $a_2 \geq \gamma^{8 \cdot 2^k} \chi^4$; coupled bound $a_8 \geq \beta^{22 \cdot 4^k} \gamma^{44k4^k + 56 \cdot 2^k}$ (formal: $\beta^{28 \cdot 4^k} \gamma^{84k4^k}$)	Thms. II.49.5, II.49.6	15
Cartesian equality (as source 14); stability with 18ε (weaker than source 14's $\tau(\varepsilon)$); spectral stability (weaker than source 14); band and profile; joint normal form (as source 14); dichotomy; vertex weights	Thms. II.51.1, II.52.2, II.53.1, Cor. II.53.2, Prop. II.53.3, Thms. II.54.2, II.55.1, II.56.1	15 (with 14)
Polynomial-loss restriction over $\mathbb{F}_p$: $\alpha^m \eta^{m-1}/(2(m+2)^m)$ above an explicit polynomial threshold, weights and vector-valued maps; order eight $\alpha^{16} \eta^{15}/(2 \cdot 18^{16})$, implying the prime case of the formal Lemma 9.3; moments γ^{r-1} without density loss (as source 04); best scalar certificate; fixed characteristic with a rank gap	Thms. F.22.1, II.61.2, Cor. II.61.3, Props. II.59.1, II.62.2, Thm. II.63.5	17 (moments also 04)
Degenerate d -arrangements: proportion $(k + C_d)/p + O(p^{-2})$, $C_8 = 2,585,444$; Lemma 15.5 with $\lambda p \geq k + 2,585,445$	Thm. II.66.8, Cor. II.66.9	18
Common sampling with lost mass at most $qd/(r+1)$; exact minimax; sharp constant $C_2 = \varphi^3 e^{-\varphi}$; shared-slice lifting; a local replacement in Lemma 16.10 (the lemma itself not proved)	Thms. II.67.4, II.68.2, II.69.4, II.70.3, Cor. II.70.8	19
<i>Part II, sources 21–27 (fifth write)</i>		

Statement	Where	Sources
Section 7 again: Lemma 7.4 with $7\delta^3 n/(3\sqrt{6})$ and the scalar optimum (as source 03); Proposition 7.3 with $(\kappa/8, 2^{15}\kappa^{-5})$ (between sources 03 and 10); order- k restriction $2^{-(34k+20)}\gamma^{12k+7}/k$, i.e. $2^{-295}\gamma^{103}$ at $k = 8$ (source 10 has $2^{-209}\gamma^{103}$)	Thms. II.74.1, II.75.1, Cor. II.77.2, Thm. F.29.1	24 (with 03, 10)
Freiman extraction for codomains $\mathbb{F}_p^d$ retaining $ B /p^{\lceil \log_p Q_k \rceil}$, independent of d ; a general moment-defect selector	Thm. II.79.5, Thm. II.73.1	24
Degenerate arrangements: degree-stratified finite bound with a degree-one class count polynomial in k ; $d = 1$: leading coefficient $(k^2 + 3k)/p$, exactly $4/p - 3/p^2$ for $k = 1$	Thms. II.83.2, II.84.1, II.85.1	22
Leading coefficient $k + H_d$ ($H_d = C_d$) over prime-power fields with codimension-one classification; residual lower bound; selection for every $d \geq 2$ with exact survival probabilities	Thms. F.30.1, II.90.1, II.92.1	25 (first order 18)
First order in every characteristic; core directions and a cubic remainder; for $d = 2$ the second coefficient $(k^2 + 21k - 6)/2$ (source 18's question answered for $d = 2$); exact count $5/q + 8/q^2 + 40/q^3 - 367/q^4 + 563/q^5 - 248/q^6$ at $k = 1$, $d = 2$, characteristic ≥ 11	Thms. II.96.3, II.97.8, F.32.3, Cor. II.99.2	27 (first order 18)
Lemma 16.10 geometry: a common base partition with width $L_m^A \geq (m/3)^A$ and exact transport to full boxes, without an inherited-coordinate hypothesis; affine gluing with $\sum_i Q_i + \sum_{i < j} Q_i Q_j$ candidates (conditional on slice covers)	Thms. II.103.6, II.103.7	23
<i>Part II, sources 28–33 (sixth write)</i> Proposition 6.1 endpoint: exact optimal near-extremal U^3 profile $1 - F \leq \mathcal{R}(\delta)$ for $\delta = 1 - \kappa < 1/120$ on every finite abelian group (equality on $\mathbb{Z}/2\mathbb{Z}$); selected graphs: graph errors $\leq \varepsilon$, infidelity $\leq (1 - \sqrt{1 - 2\varepsilon})/2$, separately sharp; weight-defect inequality and counterexamples to L^1 stability	Thms. II.107.4, II.108.1, II.109.1, Prop. II.109.2	31
The same endpoint with constants $20\varepsilon, 22\varepsilon$ (superseded by 31); soft and mixed selectors; linear order sharp	Thms. II.110.1, II.111.1, II.112.1	29
Section 7 with weights: restriction retaining $1/(kq)$ of any weight (source 10's theorem with weights); $2^{-363}c^{171}$ at order eight (source 10: $2^{-209}\gamma^{103}$); the selected derivative correlation kept through an exact restriction on $\mathbb{F}_p$ (source 02's Subsection 11.1 answered in that form)	Thms. II.116.1, II.116.2, II.117.1	31 (with 10)
Restriction over a fixed $\mathbb{F}_q$ keeping every degenerate tuple: threshold $8qS/(\alpha\eta)$, then none (source 04's Question 4 and source 17's Question 2 answered for fixed fields)	Thm. II.120.4, Cor. II.120.5	29
Section 10 by logarithmic selection: retained $\alpha p(\alpha)/20$, $O(\alpha^{-3})$ original-domain frequencies, radius $\Omega(\alpha^5)$, input defect 2^{-21} , order seven (formal: $\alpha^6/20000$, $O(\alpha^{-10})$, 2^{-43})	Thm. F.33.1, Cor. II.126.2	28 (with 07, 12)

Statement	Where	Sources
Fixed-radius escape inequality $2m(B - B \cap (B + d)) \leq B(3\rho/2) - B \leq (3^r - 1) B $, both parts sharp; Corollary 10.11 at a fixed radius $\rho/(8(3^r - 1))$; exponential dependence on the rank is necessary, even in cyclic groups	Thms. II.129.3, II.130.2, Cors. II.132.1, II.135.5	30 (with 12)
Lemma 12.3 with the linear branch $\max\{\theta^7, 4\theta - 3\}N^{32}$; dimension-free full-product repair $\delta \leq \varepsilon_s/s + 2(s-1)\varepsilon_s^2/s^2 + O(\varepsilon_s^3)$, both coefficients sharp; the sharp cubic $\varepsilon_4 \geq 4\delta - 24\delta^2 + 32\delta^3$ <i>Part II, sources 34–39 (seventh write)</i>	Cor. II.137.6, Thms. F.38.2, F.38.3	33
Lemma 15.5 by degree-weighted Fejér selection: field-size exponent $E_{k,m} = \max_t D_t/(t+1) \sim 2s/k$ (23 against 128 at $d = 8$, $k = 2$), retention $\alpha^s \eta^{s-1}/(2(s+2)^s)$, explicit threshold; a substitution at the Lemma 15.6 interface	Thm. II.145.2, Cors. II.145.3, II.148.3	35 (with 40, 22)
Sharp odd-order U^3 stability modulus $\varepsilon/4 + 3\varepsilon^2/16 + (\sqrt{2}/64)\varepsilon^{5/2} + O(\varepsilon^3)$, no fractional term when $\gcd(G , 6) = 1$; explicit local envelope	Thms. II.150.11, II.150.5	38 (with 31)
Lemma 16.10 sampling: a shared sampler keeping sampled targets, loss $\leq tq(m-r)/(m(r+1))$; closed balanced-and-capped solution of its finite minimax (for its rule; not for source 19's, [write] note); two-regime limit; derandomization; unisolvent spaces, corruption; label-preserving refinement	Thms. II.153.2, II.155.3, II.157.3, II.161.2	39 (with 19)
Sharp nonuniform transport threshold $1/(2 + (2k-2)\kappa)$ for every κ , with zero boundary loss (source 07's Question 1)	Thm. II.163.1	34 (with 07)
<i>Part II, sources 40–46 (eighth write)</i> Lemma 15.5 with one Fejér kernel: retention $\alpha^s \eta^{s-1}/(3^{s-1}(s+1)^s)$, $s = 2d2^k$, explicit field size; the single feature $r \prod_i y_i$ has exactly the intended relation space; vector targets and exact torsion losses; the barrier $a(W) \leq R(W)^{-(s-1)}$ for homogeneous positive-Fourier kernels (attained by Fejér)	Thms. F.45.1, II.166.2, II.169.2, II.172.2, Cor. II.170.1	40 (35 improves the exponent)
Lemma 15.4: core directions for every $p \geq 5$; exact core-avoidance products ($p \geq 5$ and $p = 2, 3$); the second Whitney coefficient $b_2(d)$ for every d from the nine-point grid, $b_2(8) = 3,280,234,631,541$ (second coefficient $\binom{k}{2} - 2,519,909k - b_2(8)$ at $d = 8$), stable for $p > \max(2d, 16)$; deficit of order $49^d/d$	Lem. II.175.1, Thms. II.176.2, II.176.3, II.178.1, II.179.1, II.180.1	44 (with 18, 25, 27)
The three largest torsion-free energies M_m , $M_m - 4(m-2)$, $M_m - 8(m-3)$, second level classified; coset envelope $\tau(\varepsilon)$ for $\varepsilon < 1/9$ with a canonical subgroup and all equality cases (punctured cosets); nested-hole envelope; collision equality forces one punctured coordinate	Thms. II.184.1, II.184.10, II.185.4, II.185.8, II.185.11, II.185.15	45 (with 14, 15; Pollard, Hegyvári)
<i>Part II, sources 47–53 (ninth write)</i> Lemma 9.3 made constructive: exact collision-aware inventories of energies with repeated labels, deterministic signed rounding in $O(C_s n D)$ operations, prescribed cardinality without a factor k , fixed linear systems, and a fully rational restriction with source 17's retention	Thms. F.53.1, II.190.1, II.191.2, II.192.1, II.194.1	48 (with 17)

Statement	Where	Sources
Section 13: explicit threshold $N_0(a_0)$, uniform over actual densities $a \geq a_0$, side $D(a)N^{4\gamma(a)} - 1 \geq N^{3\gamma(a)}$, mass $2^{-137}a^{704}$; full-budget witnesses with no common-step square of side two in $S \times U$ (the grid interface is not implied), repair by reselection with exponent 2γ ; the preserved all-moduli Stage 13.7 encodings false over $N = 2^k$	Thm. F.58.1, Prop. II.199.1; Thms. F.56.1, II.211.3; Thms. F.57.1, F.57.2, II.217.1, II.218.1	53, 51, 52
Two-cube degeneracy for $d = 1$: rank-two classification and census $62\binom{k}{2} + 84\binom{k}{3} + 24\binom{k}{4}$; second coefficient in every dimension with a p^{-3} remainder; $\delta_{k,1} = L_k$ exactly for $k \leq 3$ at every odd prime ($k = 3$ computer-assisted); finite fields of characteristic at least five; characteristic three exactly in every dimension; eventual Laurent polynomiality	Prop. II.222.3, Thms. II.222.4, II.223.1, II.225.2, II.224.2, Cor. II.225.5	49 (with 22)
Lemma 16.10: exact branch-cover laws for polynomial lists on progressions and product boxes, global stability, balanced words at logarithmic scales; <i>the catalogue's two-premise lemma_16_10 is false</i> at $k = \theta = \gamma = 1$ (now also formally); the full-domain endpoint product property is multiaffinity; an acyclic replacement for Corollary 16.11 conditional on I1–I7 <i>Part II, sources 54–59 (tenth write)</i>	Thms. II.228.2, II.229.1, II.233.1, Prop. II.234.1; Thm. II.235.5, Cor. II.235.11	50, 49 (with 19, 23, 39)
Proposition 6.1 endpoint: for $ f \leq 1$, $D_d(f) \leq (2/q)\varepsilon + b_d\varepsilon^2 + O_d(\varepsilon^3)$ in every degree, $b_d = (3^d - 1)/(2q^2)$ (odd order) and $(q^2 - 1)/(3q^2)$ (even order), both sharp; tangent rigidity; the exact even-order U^2 modulus; global entry by the published Eisner–Tao theorem	Thms. II.237.1, II.237.3, II.242.1, Prop. II.245.1	55 (with 31, 38)
Near-maximal uniformity: $\widehat{Q_4}(f) \geq 1 - \varepsilon$, $\varepsilon \leq 2^{-10}$, gives one scalar c with $ \widehat{\Delta_h \Delta_k f}(chk) \geq 3/4 - 4\varepsilon$ on $(1 - 332\varepsilon)p^2$ pairs; every order, finite abelian groups; a cubic inverse bound $\sqrt{1 - 6194\varepsilon} - \varepsilon$	Thms. II.247.5, II.247.7, II.247.9, Cors. II.247.10, II.247.11	57
Affine rigidity: source 11's cubic $\Psi(t) = 4t - 10t^2 + 6t^3$ for $t < 4/9$ (involution-free targets) and $t < 1/2$ (no 3-torsion either), with coset equality and endpoint rigidity; recovery for graph defect $< 2/9$ and the exact non-affine energy from $N \geq 16$ (from BCLR and source 11's cubic, as credited); relative affine repair on domains with $\ w - \alpha\ _{U^2}^4 \leq \alpha^5/1024$, error $13\varepsilon/50$, first-order coefficient $1/4$ optimal	Thms. II.249.3, II.250.4, II.251.2, Cor. II.251.3; Thms. F.61.1, F.61.3, II.260.2	59, 56 (with 11)
Section 7: exact worst-case separation ranks $\lfloor \log_2(q+1) \rfloor$ (radius $1/5$) and $\lfloor \log_3(2q+1) \rfloor$ (radius $1/7$, odd defects), sharp; weighted Freiman restrictions over arbitrary abelian targets, $ B' \geq C^{-(2k+1)\log_2(5k)} B $ from $ \Gamma - \Gamma \leq C \Gamma $	Thms. II.263.4, II.264.1, Cor. II.265.2, Thm. II.265.4	59 (with 10, 24)
Proposition 8.1: one quadratic polynomial $2^{-1}\lambda z^2 + rz$ on every translate of any P , $\sum_s \cdot \geq \beta NT/(B\sigma) \geq (27/4)\beta NT$; exact offset ceilings $1, 1/2, 1/3, 1/4$ with all equality cases; sharp stability exponents $1/2$ (order ≤ 4) and $1/4$ (order ≥ 5)	Thms. F.59.1, F.59.2, II.271.3, Prop. II.272.4	54

Statement	Where	Sources
Lemma 9.3 on every cyclic modulus and every finite abelian target: retention $(\alpha\eta/4)^{54}b^{15}$ above $b \geq (4/(\alpha\eta))^{54}$; a torsion- and collision-explicit expectation theorem; the optimal leading exponent $m \log m / (2 \log a)$ within positive Fourier filters	Thms. A, II.282.1, B, Cor. II.286.2	58 (with 04, 17, 29, 40)
Theorem 13.12 in writing: printed length $p^{2^{-\alpha}-2^{70}}$, density $(\alpha/2)^{2^{41}}$ (the printed exponent is 2^{76}), through source 40's Fejér selection and the repository's uniform square extraction, with the near-maximal endpoint for $1 - \alpha \leq 2^{-10}$	Thm. F.62.1, Thms. II.289.1, II.290.3	57 (with 40, 53)
Torsion-free energy: Schur deficit s and $n \geq 5s - 1$ force a dilate of $\{j + \lambda_j\}$, λ a partition of s ; deficit two from $n = 6$ (optimal); the third energy level $M_m - 8(m - 3)$ is attained exactly by J_m, K_m for $m \geq 7$, one more class at $m = 6$	Thms. II.292.3, II.292.4, II.292.7	57 (with 45)
<i>Part II, sources 60–69 (eleventh write)</i> Affine rigidity into elementary abelian 3-groups: $\varepsilon \geq \Psi(t)$ for all $t \leq 2/3$ with colour and target-rank remainders, coset equality for $0 < t < 2/3$; unique recovery for $\varepsilon < 1/4$ (any domain) and $\varepsilon < 4/9$ (elementary abelian 3-domains), both sharp; non-affine maximum $N^3 - 4N^2 + 10N - 6$ from $N \geq 13$, optimal	Thms. II.294.1, II.295.3, Cor. II.295.5	67 (with 11, 59)
Energy gap $m r_A(d)(m - r_A(d)) \leq m^3 - E(A)$, coefficient one, all equality cases, multilevel and weighted forms; coset existence through $\varepsilon \leq 1/4$ and uniqueness below $2/9$, both optimal; nested puncture repair with outer defect below $1/4$ and inner defect at most $1/4$	Thms. II.297.1, II.297.4, II.298.1, Cor. II.298.4, Thm. II.299.4	68 (with 14, 45)
Schur defect at most r and $n \geq 3r + 1$ force $c_i = d(i + \lambda_i)$, $p(r)$ classes; endpoint envelope $E(A) \leq M_m - 4r(m - 1 - r)$, equality only for $\{0, \dots, m\} \setminus \{r\}$; the third level from $m \geq 8$ and the fourth, $M_m - 8m + 20$, from $m \geq 11$ (two templates); torsion-free transfer	Thms. II.302.1, II.305.2, II.308.2, II.309.2	66 (with 45, 57)
Lemma 9.3 on every cyclic group, arbitrary abelian targets: exponent 53, $b \geq (4/(\alpha\eta))^{53}$, through a degree-twelve rational filter of principal cost $51.445 < 52$; uniform two-parameter bounds on the principal cost $\mathcal{I}_{m,a}$	Thms. II.312.2, II.313.1, Prop. II.312.1	68 (with 58)
Lemma 9.3 for maps $\mathbb{Z}/N\mathbb{Z} \rightarrow \mathbb{Z}/M\mathbb{Z}$ by unit-dilation averaging: retention $2^{-245}\alpha^{31}\eta^{30} B ^{15}$ above $N \geq 2^{253}(\alpha\eta)^{-31}\beta^{-15}$, and $\alpha^{16}\eta^{15}$ up to a double logarithm; the exact weighted degeneracy constant, limit $\zeta(m - 2)/\zeta(m - 1)$	Cor. II.319.2, Thm. II.314.1, Cor. II.314.2, Thm. II.317.1	62 (with 17, 58)
Exact Freiman partitions of the whole domain into at most $5mK^{(2m+1)\log_2(5m)}$ cells on arbitrary finite abelian groups; threshold-free exact order-eight restriction $2^{-2193}\alpha^{911} B $ points and $2^{-35277}\alpha^{14655} B ^{15}$ tuples; $\lfloor \log_2 m \rfloor + 1 \leq b_m \leq \log_2(5m)$	Thms. F.66.1, F.66.2, F.66.3	61 (with 10, 59)
Prime moduli: source 17's retention $\alpha^m\eta^{m-1}/(2(m+2)^m)$ under $b \geq 2(1 + \binom{m}{2})L^{m-1}$, one power of L below the intake's marked composition; an all-group certificate	Cor. II.330.5, Thm. II.330.4, Cor. II.330.3	60 (with 04, 17)

Statement	Where	Sources
Arrangements of every even order s : $\varepsilon_s \geq B_s(\delta)$, all interior equality cases ($a\mathbf{1}_C(x)\lambda(y)$, C a coset of index $n \geq 3$), every inverse coefficient optimal in succession; common-prime bracket $(s-1)/s^2 \leq c_s^{\text{prime}} \leq (3s-2)/(2s^2)$; sharp two-direction intersection modulus and biaffine repair optimal through degree $s-2$	Thms. II.334.1, II.334.3, Cor. II.335.3, Thms. II.336.3, II.337.2, II.337.7	60 (with 33)
Lemma 16.10 interpolation: certified points covered by $E_t(L_1, \dots, L_r)$ candidates (no separate direct-anchor term), exact in the fixed-cell interface by dense examples, with deletion and partition trade-offs	Thms. II.339.1, II.339.2, II.339.6	67 (with 19, 39)
<i>Part III. Cube and progression counts</i>		
No centered terms of degree 1–3; error $\sum_{j \geq 4} \binom{m}{j} \delta^{m-j} \eta^j$	Lem. III.2.1, Cor. III.2.2; Lem. III.3.1	03, 04
Exact quartic term $\delta^{m-4}(R_d \mathcal{E}(f) + T_d \mathcal{E}_{(2)}(f))$ on every finite abelian group; K_d optimal over all groups; power 4 optimal for sets; fifth-order remainder necessary; an order-three correlation	Thm. III.2.6, Section III.2	03 (odd order also 04)
$\ f\ _{U^3}^8 \geq 2\ f\ _{U^2}^8$ for real centered f on odd-order groups, equality, stability	Thm. III.3.4, Cor. III.3.5	04
Three-cube bound with sharp leading term $6\sqrt{2}\delta^4 u^4$; correction $u^{16/3}$ with optimal exponent for $\gcd(G , 6) = 1$; sharpness for sets	Thms. III.3.6, III.3.8, Props. III.3.7, III.3.9, III.3.10	04
$c_3^{\text{odd}} = 2^{-1/2}$, attained by real parts of characters (source 03's question at $d = 3$)	[write] note after Thm. III.3.4	write, from 04
Generalized von Neumann bound for invertible (admissible) slopes; telescoping error $(k-2)\alpha$ in place of $2^k\alpha$; optimal elimination order	Lem. III.1.1, Thm. III.1.2; Lem. III.4.2, Thm. III.4.3	03, 04
$T_k^*(A) \geq N^2(\delta^k - uS_{k-2}(\delta)) - \delta N$ and its thresholds	Cor. III.1.3, Cor. III.4.4	03, 04
Lifting a modular progression from a short interval	Lem. III.4.5, Rem. III.1.4	04, 03; formal
<i>Part III, sources 06–10 (second write)</i>		
Generalized von Neumann bound with two distinguished functions; progression discrepancy $\sqrt{\delta}B_k(\delta)\varepsilon^2$; sharp three-term variance factor; modular criterion	Thm. III.5.1, Cor. III.5.5, Prop. III.5.6, Cor. III.5.7	06
Cancellation through degree three; exact quartic term with order-two correction (also 03)	Thms. III.7.3, III.8.3; Thms. III.17.2, III.21.1	08, 09
Complete odd-order three-cube polynomial; Fourier lower bound $(\delta^4 + E_4)^{p/4}$; inhomogeneous weights; character-kernel survival criterion	Thm. III.9.2, Thms. III.12.1, III.13.1, Prop. III.14.1	08
Complete four- and five-vertex classifications on every finite abelian group; exact fifth order with Stirling multiplicities; sharpness through order six; transfer to sets	Thms. III.19.1, III.20.1, III.21.1, Props. III.22.1, III.22.3, Cor. III.22.5	09
The quartic lower bound with coefficient 12 is false	Rem. W1; Thm. III.10.1, Prop. III.22.2	03, 08, 09
Determinant-five configuration on six vertices; isolated ℓ -torsion for every prime ℓ	Thm. III.23.1, Prop. III.23.2	09
Torsion-sensitive mixed- L^p generalized von Neumann bound, sharp on $(\mathbb{Z}/2\mathbb{Z})^n$; density-weighted telescope with error $\delta S_k(\delta)\eta$	Thm. III.28.2, Prop. III.28.4, Thm. III.29.1, Cor. III.29.2	10

Statement	Where	Sources
Exact dimension threshold $d \leq \sum_{\ell} \lceil D_{\ell}/2 \rceil^2$ for polynomial colourings without symmetrically coloured even progressions ($D_{\ell} \leq k < p$) <i>Part III, sources 11–14 (third write)</i>	Thm. III.30.3	10
Coloop cancellation over $\mathbb{F}_p^n$; three-cube table $(1, 0, 0, 0, 12, 8, 28, 8, 1)$; exact identity with orbit moments T_5, B_1, B_2, B_3, T_7 ; spectral deficit $B_1 \leq \sigma^2 \rho^4 - \rho^8 / \sigma^2$	Lem. III.31.1, Prop. III.31.4, Thm. III.31.5, Prop. III.31.6	13 (identity also 08)
Interval discrepancy $\rho D_N(M) \leq 2\sqrt{2} \rho b^{1/4}$, exponent 1/4 optimal (Lemma 3.5)	Thm. III.32.1, Prop. III.32.2	13
Secondary coefficient $(9/4)^{1/3}$ of the three-cube extremal problem, uniformly over groups of order coprime to 6; two-harmonic rigidity; indicator sharpness	Thm. III.34.1, Thm. III.34.3	14
Source 04's Question 1 settled: the limit exists and is $(9/4)^{1/3}$	Rem. W4	write, from 14, 04
Exponent six on groups of exponent five; lower model $2^{3/4}/3$ (not claimed optimal)	Thm. III.35.2	14
Exact minimum $q^d(q^{\max(d-B, 0)} - 1)$ of symmetrically coloured progressions for polynomial colourings over $\mathbb{F}_q$; subspace equality; prime-field gap; Artin–Schreier obstruction	Thm. III.36.1, Props. III.36.6, III.36.7	14 (threshold also 10)
<i>Part III, sources 15–20 (fourth write)</i>		
Cube expansion through order five on every finite abelian group (third and fourth derivations of source 09's theorem; only 2 and 3 enter, by Smith forms); B_d sharp for sets; $2^{-1/2}$ without three-torsion	Thm. III.39.1, Prop. III.39.8; Thms. III.40.1, III.40.12	16, 18 (with 03, 08, 09)
Two marked functions with an L^2 gain for surjective homomorphisms (source 06's for slope systems); sharp finite-index comparison; sharp pair correction $\sqrt{m-1}$	Thms. III.43.3, III.44.1, Prop. III.46.1	20 (with 06, 10)
Quadratic counting $ \Lambda_k(1_A) - \delta^k \leq B_k(\delta)u^2$ with B_k strictly below source 06's $\sqrt{\delta}B_k$; supersaturation	Thm. III.45.2, Cors. III.45.4, III.47.1; [write] note there	20; write
The exponent 2 is sharp for three-term counts of indicator sets of density near 1/2 on $\mathbb{F}_{29}^{2m}$	Thm. III.48.3	20
<i>Part III, sources 21–27 (fifth write)</i>		
Exponent five: $Q_3(\delta + f) \leq \delta^8 + 6\sqrt{2}\delta^4 u^4 + (2^{3/4}/3)\delta^2 u^6 + Cu^8$ for $u \leq c\delta$, absolute constants, sharp; source 14's Question 1 answered	Thm. F.26.1; Thm. III.58.1	21, 23
Phase envelope $(2^{3/4}/3)(\cos \theta)_+^2 / (1 + \cos^2 \phi)$; rigidity; on $\mathbb{Z}/5\mathbb{Z}$ the next coefficient 20/9; indicator sharpness; a projection consequence of cube Cauchy–Schwarz	Thms. III.53.2, III.54.1, III.55.1, III.56.1, Prop. III.57.1	21 (rigidity also 23)
Polynomial colourings: exact fixed-centre pattern minimum $q^{\max(n-B, 0)}$; symmetric and monochromatic budgets in every degree; quadratic classification; no independent additive extension of the norm colouring; norm-line excess	Thms. III.61.5, III.63.1, III.64.2, III.65.3	22 (with 10, 14)
<i>Part III, sources 28–33 (sixth write)</i>		
$\ f\ _{U^d}^{2^d} \geq 2^{-m} \binom{2m}{m} \ f\ _{U^2}^{2^d}$, $m = 2^{d-2}$, for real f without self-inverse Fourier mass; $c_4^{\text{odd}} \leq (8/35)^{1/4}$ (best among sources 01–33), $c_d^{\text{odd}} \leq B_d \downarrow 1/2$, $c_d^{\text{odd}} \geq 1/3$; torsion profile $K_m(\tau)$	Thm. III.68.1, Cor. III.69.1, Prop. III.69.2, Thm. III.68.3	29

Statement	Where	Sources
$c_4^{\text{odd}} \geq 0.5587294$ by an exact two-harmonic witness (the intake's $\mathbb{Z}/7\mathbb{Z}$ certificate: ≥ 0.5587725); a single cosine is not a local extremizer for any $d \geq 4$ <i>Part III, sources 34–39 (seventh write)</i>	Cor. III.70.8, Thm. III.70.9	28; intake
$\ f\ _{U^2}^4 \leq (2/3) \ f\ _{U^4}^4$ for real centered f on odd-order groups: $c_4^{\text{odd}} \leq 2/3$ (best among sources 01–39), leading cube coefficient $200/3$; coupled bounds β_d for $5 \leq d \leq 8$ The coefficient 5 (superseded the same day by 34); $\ f\ _{U^3}^8 \geq 5S^2 - 10M + 4J$	Thm. III.71.1, Cor. III.71.9, Rem. III.71.10 Thms. III.72.1, III.72.2	34 (with 28, 29) 35
Prime-torsion onset $h(p) = D_{\text{div}}(p) + 1$; $h(7) = 7$, $h(11) = h(13) = h(17) = 8$ (source 09's question); $h(p) \sim h_{\text{cyc}}(p) \sim 2 \log p / \log \log p$; universal positivity iff inverse pairing, $h_{\text{pos}}(m) \sim 4 \log m / \log \log m$; consecutive binary determinants with cyclic Smith form (external: Alon–Vu)	Thms. III.75.2, III.76.2, III.79.1, III.80.1, III.78.1	34 (with 09)
Exponent five: exact reduction to $\mathbb{Z}/5\mathbb{Z}$ for $u/\delta \leq 10^{-5}$, so $20/9$ is the universal u^8 coefficient; u^{10} , u^{12} coefficients; cyclic reduction for every fixed odd exponent (source 21's questions)	Thms. III.84.1, III.85.1, III.86.1	38 (with 21, 23)
Exact six-term count for the linear–cubic–quintic norm colouring with an independent addition <i>Part III, sources 40–46 (eighth write)</i>	Thm. III.88.2	38 (with 22)
$\ f\ _{U^4}^{16} \geq (4795/832) \ f\ _{U^2}^{16}$ for real centered f on odd-order groups: $c_4^{\text{odd}} \leq 0.6454070$ (best among sources 01–46), leading cube coefficient 64.541 for $d = 4$ and $P_d(832/4795)^{1/4}$ for $d \geq 4$; a mean-sensitive U^3 bound; single-character functions not extremal <i>Part III, sources 47–53 (ninth write)</i>	Thm. III.90.1, Cor. III.90.2, Lem. III.90.4, Cor. III.91.2	45 (with 28, 29, 34, 35)
$\ f\ _{U^4}^{16} \geq (29/4) \ f\ _{U^2}^{16}$ (15/2 when a largest Fourier coefficient has order other than five) and $\ f\ _{U^5}^4 \geq (50/29) \ f\ _{U^2}^4$ for real centered f on odd-order groups: $c_4^{\text{odd}} \leq 0.6094183$ (best among sources 01–53), $c_5^{\text{odd}} \leq 0.58$ <i>Part III, sources 60–69 (eleventh write)</i>	Thm. III.93.1, Lem. III.93.2	49 (with 34, 35, 45)
$\ f\ _{U^2}^4 \leq (1/3 - 4/(3N^2) + 1/N^3) \ f\ _{\infty}^4$ for real centered f on every odd-order group of order N , equality on $\mathbb{Z}/N\mathbb{Z}$; $\lim_d c_d^{\text{odd}} = 1/3$, $c_d^{\text{odd}} \leq 1/3 + 10 \cdot 2^{-\lfloor d/8 \rfloor}$ for $d \geq 80$; the quartic cube coefficient $A_d = P_d c_d^{\text{odd}} \sim 6^d/24$ The same limit as $\beta_{\text{odd}} \in [1/3, 0.3464533951]$ (now $= 1/3$ by source 63) and the joint limit $1/3$ when the least prime divisor grows	Thms. III.96.4, III.97.1, III.98.5 Thms. III.99.1, III.99.2	63 (with 28, 29) 60
Gowers's Section 4: $\{0, 1, 4, 13\}$ is $(1, 5, 10)$ -separating modulo 80, so $3 \leq \beta_6 \leq \log_4 80$; an exact nine-copy circle transfer; $\rho_4(\alpha) \ll \alpha^5$, $\rho_6(\alpha) \ll \alpha^{5+9/\log_4 80}$; no perfect cyclic alphabet; the degree- $(k-1)$ dimension barrier	Thm. F.69.1, Prop. III.101.4, Thms. III.103.2, III.105.1, III.107.2	64 (with 10, 14)
Independent norm blocks avoiding symmetric $2m$ -progressions: every degree odd and below $2m$, coefficient vectors independent within each degree, for $m \leq 4$ or $q \geq (2m-1)^2$; maximum dimension rm^2 with no degree cap; eleven degree-nine certificates	Thms. III.109.1, III.109.5, III.109.6	67 (with 10, 22)

Part IV. Floor patterns

Statement	Where	Sources
$ \mathcal{F}(\mathbf{r}) = (\prod_i 2r_i)C_k$; arbitrary integer intervals and the modular count	Thm. IV.1.2; Thm. IV.6.3	01, 04
$C_k \leq \mathcal{R}(H_k, k)$, $\log_2 C_k \leq k^2 + (\log_2 e - 1)k$; $A_k \leq \mathcal{R}(k2^{k-1}, k+1)$	Thms. IV.1.4, IV.3.1; Thm. IV.6.2	01, 04
$C_k \geq 2^{\lfloor k^2/4 \rfloor}$ (01); $C_k \geq 2^{k(k-1)/2}$, $A_k \geq \prod_{j < k} (2^j + 1)$ (04; stronger for $k \geq 3$)	Thm. IV.2.1; Thm. IV.6.2	01; 04
$\log_2 C_k \sim k^2$, $\log_2 A_k \sim k^2$ (external: the classical threshold-function count)	Thms. IV.2.3, IV.3.1	01
Translation-class budgets independent of the modulus Polynomial floor entropy $(1 + o(1))n^{d+1}/d!$ (lower half external: Baldi–Vershynin)	Thm. IV.3.2; Cor. IV.6.4 Thm. IV.4.1	01, 04 01
$C_1, \dots, C_4 = 1, 2, 10, 154$; $A_1, A_2, A_3 = 2, 6, 38$ (exact rational certificates, Python verifiers)	Prop. IV.5.1; Section IV.6	01; 04
Local phase extraction at Proposition 17.7 with coefficient $2^{-k^2 - k \log_2 k - O(k)}$ in place of $2^{-2(k+1)^3}$ (not propagated to Section 18)	Thm. IV.6.7, Cor. IV.6.8; Prop. IV.3.3	04; 01
<i>Part IV, sources 06–10 (second write)</i>		
Range-carry separation for every $V \subset \mathbb{Z}_{\geq 0}^k$; modular count $\min(2r, M)^k C(V; M)$; translation classes; $\log_2 C_k \sim k^2$ uniformly in M (external: Zuev)	Thms. IV.7.3, IV.7.4, IV.7.6	06
$C_4 = 154$ certified a second time; $D_1, D_2, D_3 = 2, 6, 38$ Proposition 17.7 with coefficient $1/(8K_k^2) = 2^{-2k^2 - O(k \log k)}$	Section IV.9 Thm. IV.8.1, Cor. IV.8.2	06 (with 01, 04) 06
<i>Part IV, sources 15–20 (fourth write)</i>		
Proposition 17.7 without carry patterns: retention $(4/(3(k+2)))^{k+2}$ in the $(l+1)$ normalization, cells between $m/(3k)$ and $3km$; implies the formal statement; limit $(2/(k+2))^{k+2}$; local-scale ceiling $2^{k+1}/(k+2)!$	Thm. F.21.1, Thm. IV.13.2, Props. IV.13.4, IV.13.5	16
Exact factorial criterion for vertex polarization; functional test over $\mathbb{F}_p$; the short-cell variant (source 04’s Lemmas 8.5–8.6)	Thms. IV.12.2, IV.12.4, Prop. IV.14.2	16 (with 04)
Bits lost at $k = 8$ in one normalization: 29.07 (16), 91.07 (04, converted), 212.70 (06), 1,458 (formal)	[write] note after Cor. IV.6.8	write
<i>Part IV, sources 34–39 (seventh write)</i>		
Nonsymmetric multilinear selectors over $\mathbb{F}_p^n$ carry derivative-spectrum energy at most $1 - (1 - 1/p)^k$, sharp for $n \geq 2$; rank-profile bound p^{-r} (sharpening source 31’s noncommuting gap); lossless identity	Thms. F.41.1, F.41.2, IV.17.1, IV.18.1	36 (with 31, 29)
$\mathcal{A}_f(\sigma) = \ fe_p(-P)\ _{U^{k+1}}^{2^{k+1}}$ for $p > k+1$; noisy, dense-domain, coset, multiaffine, prime-power variants		
<i>Part IV, sources 40–46 (eighth write)</i>		
Symmetric trilinear T over $\mathbb{F}_2^n$: exact twisted cubic norm $2^{-r/2}$, $r = \text{rank}(T(a, a, b) + T(a, b, b))$, attained by one eighth-root cubic phase; integrable codimension $r/2$; agreement threshold $(1 + 2^{-r/2})/2$; seven-function correlation $\leq 2^{-r/4}$, so Tidor’s codimension $8 \log_2(1/\delta)$ becomes $2 \log_2(1/\delta)$; all-characteristic scalar polarization; trace-tensor classification	Thms. IV.27.1, IV.28.1, IV.30.2, IV.31.1	46

Statement	Where	Sources
The same theorem independently (one function, amplitude-sensitive), Boolean certification threshold $1 - 2^{1-d}$, sharp; orthogonality $\mathcal{Q}_B(g, h) \leq d_B^{-1} \ g\ _2^2 \ h\ _2^2$ on every finite abelian group; exact symmetry threshold $1 - (1 - 1/p_*(G))^k$; triangle correlation with the fourth power; quartic value $11/16$ on $\mathbb{F}_2^2$ (checked at the write, [write] note) <i>Part IV, sources 47–53 (ninth write)</i>	Thms. IV.39.4, IV.38.4, IV.35.2, IV.36.4, IV.40.1, IV.41.3	43 (with 36, 46)
Over $\mathbb{F}_2^n$: $\mathcal{M}(C_d) = 1 - (d+1)/2^d$ for the canonical tensor in every degree $d \geq 3$ and dimension $n \geq 2$ (answers source 43’s quartic question); symmetric nonintegrability thresholds $1/2$, $11/16$, $13/16$, $57/64$ in degrees 3–6, and $1 - \min\{d+1, 7\}/2^d$ in general; sharp $7/32$ non-pure trilinear support; a deficit recursion and robust certification <i>Part IV, sources 60–69 (eleventh write)</i>	Thms. F.52.1, F.52.2, IV.48.7, IV.49.1, Props. IV.50.1, IV.50.2	47 (with 43)
Over $\mathbb{F}_2^n$: the largest energy of a nonintegrable symmetric d -linear tensor is exactly $1 - (d+1)/2^d$ for every $d \geq 3$, $n \geq 2$; extremal tensors are canonical modulo integrable tensors; noncanonical cap $1 - 3d/2^{d+1}$ (65: radical-sensitive $1 - (1 - 2^{-r})d/2^{d-1}$); repair geometry, class counts $(2^n - 1)(2^{n-1} - 1)$ for $d \geq 4$, amplitude stability (65); flags $\ell \subset \Pi$, the sharp effective dimension $\text{codim } K(T) \leq \rho + \sum_{j \leq \min(\rho, d-2)} \binom{\rho}{j}$ and a fixed-radical normal form (69) Proposition 17.7: a positive face measure retains v_k/D_k with $1 \leq D_k \leq e^{2/(3(k+2))}$, so the best local coefficient is asymptotic to the ceiling $v_k = 2^{k+1}/(k+2)!$; the exact face-cover identity, positive real powers, an explicit finite theorem	Thms. F.70.1, F.70.2, IV.57.2, IV.57.4; F.74.1, F.74.2, IV.63.1	65, 69 (with 43, 46, 47)
Proposition 17.7: a positive face measure retains v_k/D_k with $1 \leq D_k \leq e^{2/(3(k+2))}$, so the best local coefficient is asymptotic to the ceiling $v_k = 2^{k+1}/(k+2)!$; the exact face-cover identity, positive real powers, an explicit finite theorem	Thms. IV.68.1, IV.68.2, IV.68.5	63 (with 16)
<i>Part V</i>		
Density and length bookkeeping for an abstract iteration	Prop. V.12.1	03
Discrete potential with optimal leading cost; explicit double-logarithmic budget	Lem. V.25.1, Thm. V.25.2	06
Step count $321(\delta_0^{-1} - 1)$; exact propagation of length losses	Props. V.42.1, V.42.2	10
<i>Open, or not certified (Section V.2)</i>		
$C_5 = 8410$ and $A_4 = 770$ (intake recount, floating-point infeasibility; <i>not certified</i>)	Item IV.2	intake

Statement	Where	Sources
The second coefficient of the $u^{16/3}$ term (between $(9/4)^{1/3}$ and 6; [6 October 2026; sources 11–14] settled: it is $(9/4)^{1/3}$, Rem. W4); c_d^{odd} for $d \geq 4$ ([6 October 2026; sources 28–33] advanced, not settled: $0.5587725 \leq c_4^{\text{odd}} \leq (8/35)^{1/4} = 0.6914416$, $c_d^{\text{odd}} \in [1/3, B_d]$, Sec. III.66; [6 October 2026; sources 34–39] upper end $2/3$ by source 34, so $0.5587725 \leq c_4^{\text{odd}} \leq 2/3$ among sources 01–39, Sec. III.71; [7 October 2026; sources 40–46] upper end 0.6454070 by source 45, so $0.5587725 \leq c_4^{\text{odd}} \leq 0.6454070$ among sources 01–46, Sec. III.89; [7 October 2026; sources 47–53] upper end 0.6094183 by source 49, so $0.5587725 \leq c_4^{\text{odd}} \leq 0.6094183$ among sources 01–53, Sec. III.92); the $k \log k$ term of $\log_2 C_k$; 46 delivered questions in all	Section V.2	all
Sources 06–10: 50 further questions and two question paragraphs (96 questions in all), merged with credits	Section V.2	06–10
Sources 11–14: 45 further questions (141 in all), merged with credits; the cutoff of source 11’s Theorem F.16.3 (a cutoff 30 via source 12’s majority lemma is a marked question); the exponent-five coefficient $2^{3/4}/3$ (numerical evidence only)	Section V.2	11–14
Sources 15–20: 74 further questions (215 in all), merged with credits; two marked questions of the intake (a combined restriction threshold; composite moduli with large prime factors); the gap between q^{-2} and q^{-1} in polynomial partitions; Lemma 16.10	Section V.2	15–20
Sources 21–27: 74 further questions (289 in all), merged with credits; the universal eighth-order coefficient in exponent five (at least $20/9$); the second degeneracy coefficient for $d \geq 3$ (reduced to one hyperplane statistic); the Section 16 closing comparison	Section V.2	21–27
Sources 28–33: 69 further questions (358 in all), merged with credits; c_4^{odd} (between 0.5587725 and 0.6914416 among sources 01–33); the full critical-norm profile beyond $\delta < 1/120$; the optimal exponential base of the fixed-radius estimate; Corollary 5.8 in degree three and beyond (source 41, placed, records a proof); one marked question of the intake (source 28 against the formal Theorem 10.13)	Section V.2	28–33
Sources 34–39: 73 further questions (431 in all), merged with credits; c_4^{odd} (between 0.5587725 and $2/3$ among sources 01–39); whether $h(p) = h_{\text{cyc}}(p)$, and the second-order support term; Lemma 5.2 beyond the first window; source 19’s profile question for its own rule; one marked question (source 35 against the formal Lemma 15.5)	Section V.2	34–39

Statement	Where	Sources
Sources 40–46: 62 further questions (493 in all), merged with credits; c_4^{odd} (between 0.5587725 and 0.6454070); the third Whitney coefficient and the third term of the degeneracy count; the largest coset-envelope range; the quartic Boolean extremum in dimension at least three; containment and rectification for Corollary 5.8; one marked question (source 40 against the formal Lemma 15.5, merged with source 35’s)	Section V.2	40–46
Sources 47–53: 57 further questions (550 in all), merged with credits; c_4^{odd} (between 0.5587725 and 0.6094183); exactness of $\delta_{k,1}$ for $k \geq 4$ and its third coefficient; the septic threshold $120/128 \leq M_{7,n}^{\text{sym}} \leq 121/128$; a faithful repair of <code>lemma_16_10</code> ; the exported Stage 13.6 length $\ell R \leq m + 1$; the simultaneous unit-step full package at Stage 13.7; near-linear constructive restriction	Section V.2	47–53
Sources 54–59: 63 further questions (613 in all), merged with credits; the cubic coefficient of the endpoint modulus and an explicit entry threshold (55); the optimal density exponent of relative repair (56); the fourth energy level and the optimal Schur threshold (57); a retention exponent $O(m)$ on every cyclic group (58); the maximal affine range and the exact cutoff between 13 and 16 (59); the best balanced-indicator coefficient (54)	Section V.2	54–59
Sources 60–69: 104 further questions (717 in all), merged with credits; the exact $\beta(q)$ on groups of fixed exponent and the optimal rate of c_d^{odd} (60, 63); the common-prime arrangement coefficient (60); the avoidance exponent b_m (61); the optimal Schur threshold, conjecturally $2r + 2$ (66); the second extremal Boolean level (65, 69); $\beta_6 = 3?$ (64); retention $O(m)$ on every cyclic group beyond positive filters (62, 68)	Section V.2	60–69

What no source claims. No source claims a new bound for $r_k(N)$; all five cite Leng, Sah and Sawhney’s $r_k(N) \ll N \exp(-(\log \log N)^{c_k})$ for $k \geq 5$ [2] as the benchmark, and say their results are local. No source claims literature priority for its formulations or an exhaustive priority search. No source compiled Lean code or claims kernel verification. Each source’s further limitations are printed in place and collected in Section V.2.

[write] 6 October 2026; sources 06–10. The same holds for sources 06–10: none claims a new bound for $r_k(N)$ (all five cite Leng–Sah–Sawhney; source 10 also Raghavan’s 2026 bound for r_3 and Green–Tao’s for r_4), none claims priority, none compiled Lean code. Sources 08 and 09 state AI assistance (provenance, Section F.4). The table rows marked “06–10” were added in the second write.

[write] 6 October 2026; sources 11–14. The same holds for sources 11–14: none claims a new bound for $r_k(N)$ (all four cite Leng–Sah–Sawhney), none claims priority (each says so explicitly; source 14 credits the classical ingredients of its colouring theorem), none compiled Lean code. Source 14 states that it was prepared by ChatGPT (provenance, Section F.4). The table rows marked “11–14” were added in the third write; in the row on the second coefficient of the $u^{16/3}$ term the gap is closed by source 14 (Remark W4).

[write] 6 October 2026; sources 15–20. The same holds for sources 15–20: none claims a new bound for $r_k(N)$ (all but source 16 cite Leng–Sah–Sawhney; source 17 also their inverse theorem), none claims priority (each says so), none compiled Lean code. Source 19’s PDF author field names ChatGPT (provenance,

Section F.4). The table rows marked “15–20” were added in the fourth write.

[write] 6 October 2026; sources 21–27. The same holds for sources 21–27: none claims a new bound for $r_k(N)$ (sources 21, 23 and 24 cite Leng–Sah–Sawhney; the others state that their results are local), none claims priority, none compiled Lean code. Sources 21, 22, 24 and 25 credit ChatGPT and source 27 calls itself AI-assisted (provenance, Section F.4). Rows of earlier writes that these sources change are not edited: the exponent-five row of source 14 (“not claimed optimal”) and the degenerate-arrangement row of source 18 stand as printed, and the rows above record what changed. The table rows marked “21–27” were added in the fifth write.

[write] 6 October 2026; sources 28–33. The same holds for sources 28–33: none claims a new bound for $r_k(N)$ (sources 28, 29 and 33 cite Leng–Sah–Sawhney; the others state that their results are local), none claims priority, none compiled Lean code. Sources 28, 29, 30, 31 and 33 credit ChatGPT (provenance, Section F.4). Rows of earlier writes that these sources change are not edited: the c_d^{odd} entry of the open row carries a dated bracket, and the rows above record what changed. The table rows marked “28–33” were added in the sixth write.

[write] 6 October 2026; sources 34–39. The same holds for sources 34–39: none claims a new bound for $r_k(N)$ (sources 35, 36 and 38 cite Leng–Sah–Sawhney; the others state that their results are local), none claims priority, none compiled Lean code. All six credit ChatGPT (provenance, Section F.4). Rows of earlier writes that these sources change are not edited: the c_d^{odd} entry of the open row carries a further dated bracket, and the rows above record what changed. The table rows marked “34–39” were added in the seventh write.

[write] 7 October 2026; sources 40–46. The same holds for sources 40–46: none claims a new bound for $r_k(N)$ (sources 40 and 46 cite Leng–Sah–Sawhney; the others state that their results are local), none claims priority, none compiled Lean code. Sources 43, 44 and 45 credit ChatGPT or AI assistance; sources 40, 41, 42 and 46 make no statement on AI assistance (provenance, Section F.4). Rows of earlier writes that these sources change are not edited: the c_d^{odd} entry of the open row carries a further dated bracket, and the rows above record what changed. The table rows marked “40–46” were added in the eighth write.

[write] 7 October 2026; sources 47–53. The same holds for sources 47–53: none claims a new bound for $r_k(N)$ (source 48 cites Leng–Sah–Sawhney; the others state that their results are local), none claims priority, none compiled Lean code. Sources 47, 48 and 49 credit ChatGPT; sources 50–53 make no statement on AI assistance (provenance, Section F.4). Rows of earlier writes that these sources change are not edited: the c_d^{odd} entry of the open row carries a further dated bracket, and the rows above record what changed. The table rows marked “47–53” were added in the ninth write.

[write] 7 October 2026; sources 54–59. The same holds for sources 54–59: none claims a new bound for $r_k(N)$ (sources 54, 57, 58 and 59 cite Leng–Sah–Sawhney; 55 and 56 state that their results are local), none claims priority, none compiled Lean code. Sources 55, 57, 58 and 59 credit ChatGPT, source 56 calls itself an “AI-generated research manuscript”, and source 54 makes no statement on AI assistance (provenance, Section F.4). Rows of earlier writes that these sources change are not edited; the rows above record what changed. The table rows marked “54–59” were added in the tenth write.

[write] 7 October 2026; sources 60–69. The same holds for sources 60–69: none claims a new bound for $r_k(N)$ (sources 60, 67 and 69 cite Leng–Sah–Sawhney; all ten say that their results are local), none claims priority, none ships Lean code. Sources 60, 61, 62, 63, 65, 66, 67 and 69 credit ChatGPT, source 68 “OpenAI ChatGPT”, and source 64 makes no statement on AI assistance (provenance, Section F.4). Rows of earlier writes that these sources change are not edited; the rows above record what changed. The table rows marked “60–69” were added in the eleventh write.

F.3 Notation across the five sources

No delivered symbol was renamed; the same letter often means different things in different sources. Each Part opens with a short table for its own sections; this is the whole dictionary. The tempting false readings are in italics.

Symbol	Meanings, by source and Part
$\mathbb{E}, \mathcal{E}, E$	$\mathbb{E}$ is always an average (expectation). $\mathcal{E}(w)$: 02, normalized additive energy of a weight on $G \times \widehat{G}$ (factor N^{-3} ; Part II). $\mathcal{E}_j(w)$: 04, $M^{-1} \sum \widehat{w} ^{2j}$ with the <i>unnormalized</i> transform, counting $2j$ -tuples (Part II). $\mathcal{E}(f) = \ f\ _{U^2}^4$ and $\mathcal{E}_{(2)}(f)$, the order-two energy: 03 (Part III). $\mathcal{E}(\lambda)$: 04, the additive energy of $\lambda_r = a_r ^2$ (Part III). E : 05, captured Fourier energy $\sum c_i ^2$ (Part I); 01, an exceptional set (Part I); 02 and 03, the additive energy $E(\Gamma)$, $E(A)$ of a graph or a set (Part II); 04, E and E_j in Lemma IV.6.6 (Part IV). E_{osc} : 03, an oscillation error. <i>None of these is a normalization of another.</i>
Q	03, 04: Q_d the normalized cube count (the same quantity, Part III). 01: Q_i box counts (Part I). 05: $Q_i(L)$ box counts and Q in Lemma I.16.1 (Part I). 02: Q and $Q_{a,b}(p, t)$, good-mass certificates; Q_i input progressions (Part I). 04: Q the kernel ratio sum (R14), Q_K (R32), $Q_k = k2^{k-1}$, intervals Q_i (Parts II, IV); questions Q1–Q12 and displays (Q1), (Q2) (Part V). 03: Q in Lemma I.15.1.
$R, \mathcal{R}$	03: R_d the number of ordinary parallelograms of $\{0, 1\}^d$, $\mathcal{R}_d(F)$ a remainder (Part III). 04: P_d the same number; R_d, R_6, R_7 remainders; R_j counts of respected arrangements (Part II); $R_d(H) = \sum_{j \leq d} \binom{H}{j}$ (Part IV); R_i cells. 01: $\mathcal{R}(h, d) = \sum_{j \leq d} \binom{h}{j}$ (Part IV), <i>so 01's $\mathcal{R}(H, d)$ is 04's $R_d(H)$, arguments reversed.</i> 02: $R = \mathbb{E} K_u(s) ^2$ (Part II), R_0 the partition threshold, R_j refined cells (Part I). 05: R an oscillation budget (Part I). 03: R the number of steps (Part V). <i>04's remainder R_d is not 03's count R_d.</i>
$T, \mathcal{T}$	02: $\mathcal{T}_w(f, g)$ and T , weighted derivative correlations (Part II). 03: T_d binary tetrahedra (Part III); $T = \mathbb{E}X_+$, T_0 (Part I). 04: $T_d(B)$, $T_d^\phi(B)$ tuple counts (Part II); $T(s, t)$ Fourier cube coefficients, $T_k^*(A)$, $T_i(a)$ (Parts III, IV). 01: T_n , $T(n, d)$ threshold-function counts (Part IV); $T = \prod_i Q_i$ (Part I).
κ	01: $\eta/(2\delta(1 - \delta))$ (Part I). 02: $\mathcal{U}(f)/m(f)^4$ (Part II). 03: $7/(3\sqrt{6})$ and κ_* (Part II). 04: $\kappa_d = 1/(2(400d)^d)$ (Part II). 05: a Gram-matrix bound (Part I).
ρ	02: half the chordal diameter of a phase set (Part I). 03: a normalized Fourier coefficient (Part I); $c/(2 - c)$ (Part II). 04: m/M (Part II); the largest Fourier modulus (Part III). 01: $\max\{0, \lambda_* - \beta_E\}$ (Corollary I.2.4).
s	01: $\sin \theta$ for an arc of half-angle θ . 05: a positive mass, and $s(\varepsilon)$. 02: a weighted absolute mean. 03: $\sqrt{1 - \gamma/\varepsilon}$ (Part II). <i>01's s is 02's ρ, not 05's s.</i>
C_k, C_d, C	01: C_k fractional carry patterns (= 04's F_k , Part IV). 04: $C_d = 1 + \binom{2d}{2}$ (Part II). 02: C a local correlation, C_0 (Part I).
B, B_k	01: $B_k = (2(k+1))^k \mathcal{R}(k2^{k-1}, k+1)$ and $B_{\delta,s}(a)$. 04: $B_k = (2k)^k R_{k+1}(k2^{k-1})$, $B_5(f)$, the set $B \subset \mathbb{F}_p$. 02: B_0, B_ρ , a set B . 05: $B = \mathbb{E} f $, a width budget B . 03: $B(S)$ (= 04's $D(S)$). <i>01's and 04's B_k are different numbers.</i>
L	02: the total absolute mass $\sum f $. 03: $L = D $ (Part II), a length (Part I). 04: kernel bandwidth (Part II), number of labels (Part IV), cube forms L_v (Part III, also 03). 01, 05: cell lengths.
M, m, q	M : 01 $M_\delta = 2\delta(1 - \delta)$ and a modulus; 02, 03 a number of cells; 04 $ H $, $\sum \lambda^4$, $M_\delta(u)$, a modulus; 05 the modulus of Lemma 2.3. m : 03 2^d (cubes), $ V $ (DRC), $ A $ (energy); 04 $ B $, $ A $, $2h+1$; 01 anchor count, monomial count; 05 the number of phases. q : 04 2^d (Part III), kernel degree (Part II), $\lceil h/k \rceil$ (Part IV); 01, 03, 05 a common difference; 01 q_j cell densities; 02 q_t good mass.

Symbol	Meanings, by source and Part
α, η, u	Gowers's “ α -uniform of degree $d - 1$ ” means $\ f\ _{U^d}^{2^d} \leq \alpha$: α is a norm power, not a norm (03, 04). Elsewhere α is a correlation (02, 05, Gowers's Corollaries 2.4-2.5) or a retained proportion (04, Part II). η : 01 a normalized correlation; 03 a U^d norm (Part III) or a cell discrepancy bound; 04 the unrespected proportion (Part II). u : 04 a U^d norm.
arcs	01: an arc of r turns, half-angle $\theta = \pi r$, $s = \sin \theta$. 05: an arc of width ε turns, disk radius $\sin(\pi\varepsilon)$. 02: chordal diameter 2ρ , so $\rho = \sin \theta$. 03: chordal error ε . <i>A width in turns is not an angle.</i>
Fourier	02, 03, 04 (outside its Section II.8), 05: normalized, $\widehat{f}(\chi) = \mathbb{E}_x f(x) \overline{\chi(x)}$. 04 Section II.8: unnormalized. 01: normalized correlations $\mathbb{E}(h - \delta)e(\xi n)$, and the unnormalized cyclic transform in Corollary I.6.2. The Lean development: unnormalized (<code>ZMod.dft</code>), so its hypothesis $ \widehat{1_A}(r) \geq \alpha N$ is a normalized coefficient at least α .
Gowers norms	02: $\mathcal{U}(f) = \ f\ _{U^3}^8$, normalized; its repository counterpart is $N^4 \mathcal{U}(f)$, and the degree parameter of <code>UniformOfDegree</code> is 2. 03, 04: $\ f\ _{U^d}$ normalized; Lean's <code>gowersNorm</code> $d f = N^{(d+1)/2^d} \ f\ _{U^d}$.
numbering	“Lemma 5.15”, “Corollary 2.4”, “Section 17” and other bare numbers are Gowers's. This report's numbers carry a Part numeral (I–V) or F. Source 04's display tags (Mj), (Rj), (Aj) are not this report's equation numbers; <i>(A1) is not Appendix A, and (R3) is not a remainder.</i>
<i>Sources 06–10 (added in the second write; the Part heads have the full rows)</i>	
η, δ, γ	06: η a density threshold (Part I). 07: δ, δ_d failure probabilities of directions, <i>not densities</i> (Part II). 10: η a correlation (Part I) or a U^{k-1} norm (Part III); γ a retained correlation fraction (Part I) or a graph-energy proportion (Part II).
$\mathcal{C}, \mathcal{C}_d, \mathcal{C}_d$	07: $\mathcal{C}(\mathbf{d})$ the cost of a word (Part II). 08: $\mathcal{C}_d(F)$, 09: $\mathcal{C}_d(F)$, the normalized cube count, = 03's and 04's Q_d (Part III).
$Q_d, R_d, P_d, T_d, F_d, A_d, K_d, S_d$	08's $R_d = 09$'s $P_d = 03$'s $R_d = 04$'s P_d ; T_d the same in 03, 08, 09; 08's $F_d = 09$'s $A_d = 03$'s K_d . 08's $Q_d(f)$ (the quartic coefficient), 09's Q_d (star count) and 08's $K_d = \sum_{r \geq 4} \binom{p}{r}$ are not 03's or 04's Q_d, K_d ; 09's S_d counts ternary supports.
κ	07: a fibre-size ratio. 10: $\kappa_G(c) = \{x : cx = 0\} $. 08: $\kappa_d^{\text{odd}} = 03$'s c_d^{odd} .
B_k, c_k, h_k, H_k	06: $B_k(\delta) = \sum_t (t+1)\delta^t$ (Part III), $K_k = 01$'s B_k and $c_k = 1/(8K_k^2)$ (Part IV), $h_k = (k-2)2^{k-1} + 1 = H_k$ of 01, 04. 10: H_k, h_k a vertical obstruction fibre and its size (Part II). <i>06's c_k is not 04's c_k.</i>
M, v, S	06: M the number of cells, v the cell variation. 10: M the largest cell density, $v = \delta(1 - \delta)$, $S_{\mathcal{P}}$ the first moment (= 06's v for indicators).
<i>Sources 11–14 (added in the third write; the Part heads have the full rows)</i>	
κ	11: $\kappa \in \{0, 1\}$, whether the target has an involution (Part II). 12: a bound on the translation loss b_B (Part II). 14: the constant $2^{5/4}/12$ (Part III). <i>None of these is any earlier κ.</i>
$\mathcal{E}, E$	11: $\mathcal{E}(f)$, the ordered graph energy of $f : G \rightarrow W$, and $E(S)$, the set energy of the error support; 14: $E(A)$ (= 11's $E(S) = 03$'s $E(A)$) and $\mathcal{E}(\lambda)$ (= 04's, Part III). 12: E_t is a set of t -edges, not an energy.
χ, τ	12: $\chi(\eta) = (1 - \sqrt{1 - 2\eta})/2$, the majority function, <i>not a character</i> ; 14: $\tau(\varepsilon) = (1 - \sqrt{1 - 4\varepsilon})/2 = \chi(2\varepsilon)$ (Part II).
ρ, ϱ	12: ρ an outer Bohr radius, ϱ the component parameter of Lemma 10.5. 13: $\rho = \ f\ _{U^2}$ (Part III), and ρ_0 a chord bound (Part I). 14: $\rho = \widehat{f}(\xi) $, half of source 04's amplitude a (Part III). 11: ρ a test failure probability.
$\varepsilon, \eta, \theta$	11: $\varepsilon(f) = 1 - \mathcal{E}(f)/N^3$ and η a near-extremal slack. 12: ε a missing density, η a collision defect or 2^{-43} , $\theta = 10\eta^{1/5}$. 14: ε an energy defect, $\eta = S_0 - S$ a norm deficit.
D, d	11: D_κ the inverse cubic, $d(f)$ the distance to affine maps. 12: D the canonical difference map, $d = \Gamma $ the rank. 13: D the cell discrepancy (= 01's D) and $D_N(M)$ the interval constant; d a common difference or a cube dimension. 14: D a Fourier tail mass.

Symbol	Meanings, by source and Part
T, B	11: T non-monochromatic triples. 13: T_5, T_7, B_1 – B_3 cube moments (13's $T_5 = 04$'s B_5). 14: $T(s, t)$ Fourier cube coefficients (= 04's), B a degree budget (= 10's B), B_r cross-sections.
A, C, J, M	11: colour statistics of the error values (counts, <i>not sets</i>). 14: A a set, C a coset, M a common cube mass or $\sum a_r ^8$.
<i>Sources 15–20 (added in the fourth write; the Part heads have the full rows)</i>	
C, C_d	15: $C = R_1$, a collision count (<i>not 14's coset C</i>). 18: $C_d = (T_{2d} - 2\binom{2d}{d} + 1)/2$, a degeneracy constant (<i>not 01's or 06's carry counts C_k</i>). 19: C_d a Poisson sampling constant, $C_2 = \varphi^3 e^{-\varphi}$.
Q, Q_d	15: $Q(A)$ parallel axis-cube pairs (= 14's $Q(A)$). 16: $\mathcal{Q}_d$ the cube functional (= 03's and 04's Q_d), Q a cell. 18: Q_d five-circuit counts (= 09's $Q_d = 16$'s F_d) and $Q_d(p)$ a probability bound. 20: $Q_d(f, g)$ a pair correlation.
$B, B_k, \mathcal{B}_k$	16: $\mathcal{B}_k(g; \mathcal{Q}, U)$ the partition quotient (<i>not 01's or 04's B_k</i>). 18: B_d the all-group quartic count. 20: $B_k(\delta)$ (<i>not 06's B_k</i> : 20's is at most $\sqrt{\delta}$ times 06's).
R, T, S	15: R_d arrangement counts (= 04's), T the vertical group (14 writes R). 17: T_m with m variables (17's $T_{16} = 04$'s T_8), $S_{m,L}$ (= 04's Q). 18: R_d (= 03's R_d), T_n central trinomial coefficients. 19: $S = \sum_i d_i(d_i + 1)/2$. 20: $R_m(\delta)$, $S_{k-2}(\delta)$ (= 03's).
χ, σ, ρ_3	15: $\chi = C/(M^3 n)$, a normalized collision count (<i>not a character, not 12's $\chi(\eta)$</i>), ρ_3 (= 04's σ). 19: σ a lost mass. 20: $\sigma = \sqrt{\delta(1 - \delta)}$.
q, m, L, d	16: q the number of affine frequencies (= 13's q ; <i>not 04's $\lceil h/k \rceil$</i>); $m = 2h + 1$, or 2^d in its cube section. 17: m the number of variables, $L = \lceil (m + 1)/(\alpha\eta) \rceil$ (04's $\lceil 4/(\alpha\eta) \rceil$), q a filter mean. 18: $q = 2d$ (Section 15) or the number of phases; d an arrangement order or a cube dimension. 19: q classes per row, d an interpolation order.
A, E, D, J	18: A_n ternary classes (<i>not 01's or 04's A_k</i>), E_k a selection exponent, E_2, E_3 torsion moments, $D = (2d + 1)k + 2d - 1$, $J_{k,d}$. 19: $D = \sum_i d_i$. 16: D (= 04's B_5).
κ, λ, φ	15: κ a balance constant (= 04's K). 18: $\kappa > 10.5$ (Lau), λ a selection density. 19: κ_j recurrence constants, λ a Poisson mean, φ the golden ratio (<i>not a map</i>).
degenerate	17: L -degenerate tuples (Section 9) <i>are not</i> 18's degenerate arrangements (Section 15).
<i>Sources 21–27 (added in the fifth write; the Part heads have the full rows)</i>	
C_5, C_d, H_d, L_d	21, 23: $C_5 = 2^{3/4}/3$, the exponent-five coefficient (= 14's). 22: $C_d = 18$'s C_d ; 25: $H_d = 18$'s C_d ; 27: $L_d = 18$'s C_d (<i>the same number $(T_{2d} - 2\binom{2d}{d} + 1)/2$ under three names; not 01's or 06's carry counts C_k</i>). 24: C_t scalar constants ($C_3 = 0.9097\dots$, $C_2 < 0$) and C a doubling constant. 21: C_6, C_7 six- and seven-vertex sums (<i>not 23's C_6, one of three six-vertex forms</i>).
$H, J, \mathbf{J}$	21: H a cubic statistic (= 23's A_6). 22: $J_{k,d}$ balanced degree-one classes (<i>not 18's $J_{k,d}$, all nonparity classes, which 22 writes J^{old}</i>), H_t a Hamming ball, H a hyperplane; 23: H the largest allowed length; 25: $\mathcal{H}_d$ the hyperplane family.
B, B_5	21, 23: B_5 the quintic Fourier sum (= 04's B_5). 22: $B(\Pi, D)$, B_{sym} , B_{mono} pattern budgets (= 14's B). 23: B the Dirichlet bound; 27: $B_\ell(p)$; 18: B_d .
q, L, Q	22: $q = 2d$ in Sections 2–5, the field size in Sections 6–10. 23: q the number of phases (13's sense), $L = \min \mathcal{S}$. 26: L a progression length, Q the output progression. 24: Q_k the vertical defect fibre size (= 10's h_k). 27: $Q_\ell(p)$ profile classes; 18: $Q_d(p)$.
S, R, U	27: $S_{k,d}$ core directions, $R_{k,d}$ non-column-parity directions ($R_{2,2} = 2,585,419$ <i>is not</i> $H_8 = 2,585,444$), U_ℓ unbalanced classes. 25: $U_{k,d}$ a second-order constant, R_n a survival ratio. 22: U_t a cumulative profile. 24: S a set with small difference set.
$\kappa, \beta_5, \eta, \theta$	23: $\kappa = 2^{1/4}/6 = 21$'s β_5 (<i>not 18's or 19's κ</i>). 24: $\kappa = E(A)/m^3$ (= 10's γ , 03's c). 21, 23: $\eta = S_0 - S$; 22: η a certificate; 23: η an exception budget. 21: $\theta = 2\alpha - \beta$; 23: $\theta = \arg a_\xi$.
$\mu, \tau, \alpha, \delta$	26: μ the normalized balanced mass, τ the cover density, α a summed correlation, δ the density. 24: μ an index measure, τ a popularity threshold, δ Lemma 7.4's density. 25: τ a tolerated fraction.
<i>Sources 28–33 (added in the sixth write; the Part heads have the full rows)</i>	

Symbol	Meanings, by source and Part
$c_d^{\text{odd}}, C_m, B_d$	28, 29: $c_d^{\text{odd}} = 03's\ c_d^{\text{odd}} = 08's\ \kappa_d^{\text{odd}}$. 29: $C_m = 2^{-m} \binom{2m}{m}$ a moment coefficient (<i>not 01's or 06's carry counts C_k, 18's C_d or 21's C_5</i>), $B_d = 2 \binom{2m}{m}^{-1/m}$ an upper bound for c_d^{odd} (<i>not 18's B_d</i>), $K_m(\tau)$ a torsion profile.
κ, δ	31: $\kappa = \ f\ _{U^3}^8 / \ f\ _2^8$, $\delta = 1 - \kappa$ (<i>not 07's fibre ratio or 10's κ_G</i>). 30: κ a normalized translation loss (= 12's κ). 33: δ the distance to the model class. 32: δ the density.
B, T, H, q	28: Bohr(Γ, u) Bohr sets, T an anchored energy. 30: $B(\chi, \rho)$ rectangular Bohr sets, $b_B(d)$ the translation defect (<i>a count, where 12's b_B is a fraction</i>). 33: B a biadditive map. 29, 31: T_ϕ the same derivative score. 29, 31: H an isotropic subgroup; 28: H the target group; 32: H a block length. 29: q the field size; 31: $q = V_k $; 32: q a Dirichlet denominator.
<i>Sources 34–39 (added in the seventh write; the Part heads have the full rows)</i>	
$c_d^{\text{odd}}, C_m, H_m(p)$	34, 35 = 28, 29, 03: c_d^{odd} . 34: $C_m = 2^{-m} \binom{2m}{m}$ (= 29's), $H_m(p)$ a moment profile of the largest pair mass p (<i>not a prime</i>), $K = 222^{1/4}/2$. 35: $C(S)$ the conductor of a set of lengths, $C_{k,m,\beta}$ a constant (<i>not 21's C_5, 18's C_d or 29's C_m</i>).
$h(p), D, W$	34: $h(p)$, h_{cyc} , h_{pos} (= 09's $h(\ell)$), $D_{\text{div}}(p)$. 35: D_t a dimension, D a divisor product, $W = \sum_i d_i(d_i + 1)/2$ (= 19's S) in its Section 7 but $W = \sigma^\perp$ in its Sections 2–5; 37: D a Weyl exponent; 39: W a total row weight; 19: $D = \sum_i d_i$.
$M, L, t, \kappa, \rho, \lambda$	35: $M = (m + 1)k + m - 1$ (Sections 2–5) and $M = \sum \lambda^4$ (Section 6); 38: $M = \sum a ^8$, $\mathcal{M}_{\delta,G}(u)$; 39: $\mathfrak{M}_{m,q,r,t}$; 37: M a gap half-width. L : 34 a pair statistic, 35 a Fejér bandwidth, 39 a list size. t : 34 a two-pair parameter, 37 $ A $, 39 the anchor threshold (= 19's d). κ : 34 the fibre ratio (= 07's), 38 $\ v\ _{U^3}^8$ (= 31's), 39 κ_t . ρ : 37 $\rho_R(a)$, 34 a pair amplitude. λ : 35, 38 $ a_r ^2$; 37 $(1 - c)/S$; 39 λ_t .
<i>Sources 40–46 (added in the eighth write; the Part heads have the full rows)</i>	
s, M, D, K	40: $s = 2d2^k$ (= 35's s), $M = (2d + 1)k + 2d - 1$ (= 35's M , 44's D), $E_k = s2^{s-1}$ (Gowers's exponent). 41: $s = k - 1$ differencing shifts (<i>not 40's s</i>), D_k the density exponent, $K_k = (k!)^2 2^{(k+1)^2}$ (= the formal polynomialPartitionConstant), A_k, B_k recurrence constants (<i>not 37's A_k</i>), c_k, e_k budgets, $\mathcal{P}_k = k!2^{k(k+1)/2}$, M the number of parents. 42: $M_L(A), H_L(A)$, m the largest divisor of N below L , $q = N/m$ (<i>not a field size</i>), K a subgroup. 44: D the parameter count, H_d (= 18's C_d), $b_{2,d}$ (= 27's $b_{2,d,p}$), $m = 2d$.
M_m, τ, η, κ	45: $M_m = (2m^3 + m)/3$ the torsion-free maximum, $m = A $, H_m, J_m sets (<i>not 44's H_d</i>), $\tau(s) = (1 - \sqrt{1 - 4s})/2$ (= 14's), $\eta = \varepsilon - \delta(1 - \delta)$ (<i>not 40's accuracy or 42's collision probability</i>), $Q_d = \ f\ _{U^d}^{2^d}$, $S = Q_2$ (= 29's S), $\kappa = (832/4795)^{1/4}$ (<i>not 34's, 38's or 31's κ</i>). 41, 42: $\tau = U /N$.
C_T, Ω_T, r, m	46: $C_T(a, b) = T(a, a, b) + T(a, b, b)$, $r = \text{rank } C_T$, $m = r/2$, Λ_T, Γ_T (<i>not 40's Γ_s</i>), $Q(a) = T(a, a, a)$. 43: Ω_T (= 46's C_T), $\mathcal{E}_T(f)$ the tested energy, $\mathcal{M}(T)$ its maximum, $\mathcal{Q}_B, d_B, \text{nCSM}^d, C_d$ a quartic family (<i>not 18's C_d</i>).
<i>Sources 47–53 (added in the ninth write; the Part heads have the full rows)</i>	
p, q, M, E, W, C, H	49: p a prime in Sections 2–6 and 8, but the largest inverse-pair mass in Section 7 (it says so); $q = p^m$; M the principal dual coefficient (<i>not 44's or 45's M_m</i>), E its dual tail, W a norm bound, $C = (2M)^{16}/(16W^{15})$, $H_m(p)$ (= 34's). 50: p the prime, q a list length and $q(\cdot)$ a control function, $E = E_2 = 2^{1024}$, $Q = 2^E$, R an alphabet size, $c = 1/Q$ (<i>not 49's</i>). 48: q a new probability, $H_p(t)$ a response series, R a respected count, T a total count, $C_s = \binom{s+2}{2}^2$.
$K, Q, \gamma, \beta, c, f, g, T, R$	51–53: $K(a) = 2^{114} a^{-320}$, $Q(a) = 2^{220} a^{-221}$ (= the formal section13Q), $c = \zeta_{13}/2$, $\beta = 2^{-13}Q$, $f = 2^{-100} a^{448}$, $g = 2^{-284} a^{1408}$, $\gamma = \beta f g / 4$ (= section13SquareExponent ; <i>not 50's γ</i>), R the Stage 13.6 progression (<i>not 50's alphabet size</i>), T the strip or a fixed exponent. 47: C_d the canonical tensor (<i>not 18's or 44's C_d</i>), $\mathcal{M}(T)$, $\mathcal{E}_T(f)$ (= 43's), $\mathcal{I}_d(V)$ (= 43's nCSM^d), $b_d = 1 - (d + 1)/2^d$.
<i>Sources 54–59 (added in the tenth write; the Part II head has the full rows)</i>	

Symbol	Meanings, by source and Part
Q, D, Ψ, P	55: $Q_d(r)$ a quartic circuit sum, $D_d(f)$ the squared distance to phase polynomials; 57: $Q_d(f) = \ f\ _{U^d}^{2^d}$ (= 45's), $D(A)$ the energy defect, P, Q progressions; 56: $Q(w)$ the weighted quadruple mass, $\mathcal{D}_w(\phi)$ the defect; 58: $Q = P_2$ an even filter weight, P the principal weight, D the filter degree; 59: $\Psi(t) = 4t - 10t^2 + 6t^3$ (= 11's Ψ_0), $D(u)$ its inverse (<i>not</i> 55's D_d), P zero-sum pairs; 54: P the translate set, $D_h(r)$ Gowers's coefficient, $\mathcal{D}_\eta$ the spectral overlap.
$\varepsilon, t, a, c_m, C_d$	55: $\varepsilon = 1 - \ f\ _{U^d}^{2^d}$, a an amplitude deficit; 56: ε the relative test error, t the relative distance; 57: $a = \alpha/2$, C_s correction constants; 59: $\varepsilon(f)$ the graph defect, t the error density, a a colour count; 54: c_m offset ceilings; 55: C_d (= 04's P_d), A_d, J_d, B_d cube counts; 58: c_n filter coefficients (<i>none of these is another source's</i>).
<i>Sources 60–69 (added in the eleventh write; the Part heads have the full rows)</i>	
$c_d^{\text{odd}}, \beta, b(G)$	60, 63: c_d^{odd} as sources 28 and 29 (the ratio of fourth powers, not a 2^d -th-power coefficient); 60: $\beta_{\text{odd}}, \beta(q)$ bounded Fourier extrema; 63: $b(G)$ the same extremum on one group; 64: β_6 a growth exponent and $b_6 = \log_4 80$; 62: $\beta = B /N$ (<i>four different β's</i>).
ρ, B, K	64: $\rho_k(\alpha)$ Deng–Tidor–Zhao's progression constant; 60: ρ_s the inverse branch of B_s ; 69: $\rho(T)$ a radical codimension (= 65's $r(T)$); 60: $B_s(\delta)$ the arrangement polynomial; 69: $B_d(r)$ a dimension bound; 65: K_T the contraction radical (= 69's $R(T)$, <i>not</i> 69's $K(T)$); 61: K a difference ratio; 62: K_m a constant.
E, ε, t	66: $E(A)$, 68: $E(A)$ the additive energy (= 45's); 67: E_t an elementary symmetric polynomial, $\mathcal{E}(f)$ the graph energy (= 59's); 65, 69: $\mathcal{E}_T$ the selected energy; 68: $\varepsilon = 1 - E(A)/m^3$; 67: $\varepsilon(f)$ the graph defect, t the error density (= 59's), and in its Section 5 t a degree bound; 60: ε_s a failure probability.

F.4 Provenance and merge decisions

Archives. The five archives arrived in commit 66f24b0d0 (“New research reports”, 6 October 2026) and were placed in this directory by aefb0b44a (“Place batch 115 (1/2)”, batch 115 of ProveIt's incoming-reports intake), which also retired them from docs/incoming/. They survive in the arrival commit (git show 66f24b0d0:docs/incoming/<archive>).

No.	Archive	Bytes	Files	Pin	Manuscript
01	Gowers_Quantitative_Refinements(1).zip	541,437	15	3a25020a6	1,980 lines, 31 pp.
02	Gowers_Szemerédi_Refinements.zip	492,008	7	(66c02f8a1)	823 lines, 20 pp.
03	gowers_local_refinements.zip	408,961	10	3a25020a6	2,392 lines, 36 pp.
04	gowers_quantitative_refinements.zip	788,239	14	327773149	2,987 lines, 40 pp.
05	gowers_strengthening.zip	399,349	8	327773149	847 lines, 18 pp.

The pins are three commits of 6 October 2026, in the order 3a25020a6 → 327773149 → 66c02f8a1 → ... → HEAD. Every repository blob that a source cites was checked at intake against its pin and found correct.

Authorship and disclosures. Sources 02, 03 and 04 are “prepared for Vladimir Reshetnikov” (title block and PDF author field), source 05 is a “research note prepared for Vladimir Reshetnikov” with an empty PDF author field. *Source 01 credits ChatGPT*: its author line reads “Research report prepared for Vladimir Reshetnikov / Mathematical development and computational checks with

ChatGPT”, and its PDF author field reads “Research report prepared with ChatGPT”. No source names a human author. Source 03 mentions “independent mathematical reviews” that are not part of its delivery (note in Section V.13).

What each source contributed. 01: the sharp weighted density-tail theory with a phase parameter, simultaneous partitions, the strongest single-coefficient Fourier increment, the scale barrier (Part I), and the carry-pattern count with $\log_2 C_k \sim k^2$ and $C_4 = 154$ (Part IV). 02: the mixed weighted extension of Proposition 6.1 (Part II), sharp real phase freezing, the general sharp tail bound, the Lemma 5.15 replacement and the refinement theorem under $\text{PP}(K, R_0)$ (Part I). 03: the selection theorem and a Fourier increment (Part I), dependent random choice and an elementary Proposition 7.3 (Part II), progression counting and the exact quartic cube term on every group (Part III), iteration bookkeeping (Part V). 04: moments and polynomial restriction (Part II), the sharp odd-order cube constants (Part III), floor patterns with the stronger lower bound and the Proposition 17.7 coefficient (Part IV). 05: the sharp constant $1/3$, Corollary 2.4 with 36π , a density-sensitive increment with popularity, simultaneous localization with its optimal exponent, spectral extraction (Part I).

Overlaps. The five are independent manuscripts: no two share more than 0.8% of their text (intake measurement, 8-word shingles), and sources 01 and 04 are not editions of one another (different titles, pins and files; 0.23% shared text). The results proved by several sources are listed in Section F.2 with all their sources.

Where the merge had to choose.

- *Base.* Source 04 is the base: it was the largest and broadest at intake; its manuscript was staged as `article.tex` and its delivery README as `README.md`, both now replaced by this report and its README.
- *Thematic Parts*, by the intake’s recommendation and Vladimir’s direction, since every source spans two to four themes. Each delivered section is kept whole in exactly one Part; each source’s abstract and Section 1 are in the front matter. Sources’ conventions sections go with their first use: source 03’s Section 2 opens its Part I chapter, source 02’s Section 2 and source 04’s Section 2 open their Part II chapters. Source 02’s Appendix B, a rational example for its tail theorem, is in Part I with that theorem; its Appendix A is in Part V.
- *Results proved by several sources* are printed in each source’s own form, because the forms and hypotheses differ (general $[-a, b]$ against $a = \delta$; strict against non-strict events; all groups against odd-order groups), and because the proofs are independent, mostly by different routes. A dated note at each later print names the first print and credits every source; where two proofs follow the same argument, the note says so. No proof was replaced by a pointer.
- *The sources’ question sections* are printed as delivered, and the write’s Section V.2 deduplicates them with credits.

[write] **6 October 2026; sources 06–10.** *Sources 06–10.* The five archives arrived in commit 62e21161f (6 October 2026) and were placed in this directory by ad37962c4 (“Place batch 115 (2/2)”), which retired them from `docs/incoming/`; they survive in the arrival commit (`git show 62e21161f:docs/incoming/<archive>`). Not staged: the five PDFs and the five manuscripts and delivery READMEs, whose text is printed here.

No.	Archive	Bytes	Files	Pin	Manuscript
06	Gowers_Refinements (1).zip	647,760	19	7d99e5dee	1,958 lines, 27 pp.
07	Gowers_Transport_Rigidity.zip	359,160	9	ae0e89574	1,472 lines, 21 pp.
08	gowers_cube_refinements.zip	541,854	8	(blobs)	1,677 lines, 25 pp.
09	gowers_cubical_expansions.zip	512,022	10	7d99e5dee	1,794 lines, 27 pp.
10	gowers_refinements.zip	664,416	15	7d99e5dee	2,443 lines, 34 pp.

The pins precede HEAD; ae0e89574 is a descendant of 7d99e5dee. Source 08 records blob identifiers, not a commit; they are those of 7d99e5dee. All blob identifiers the sources cite were checked at intake.

Authorship and disclosures. Source 06: “Research report prepared for Vladimir Reshetnikov” (title block and PDF author field). Source 07: “Research draft” (title page, running head, PDF author field). *Source 08* is “Prepared for Vladimir Reshetnikov”, PDF author field “Research memorandum prepared for Vladimir Reshetnikov”, and its title page calls it “This AI-assisted research draft”. *Source 09* credits ChatGPT: its author line reads “Research manuscript prepared with ChatGPT”, and so does its PDF author field. Source 10: “A mathematical research report”, PDF author field “Research report”. No source names a human author. Source 10 mentions an “independent review” that is not part of its delivery (note in Section V.43).

What each contributed. 06: the exact fixed- M extraction (Part I), a two-target generalized von Neumann bound and quadratic progression discrepancy (Part III), carry counts for general families with entropy uniform in the modulus, a second certificate of $C_4 = 154$, and a Proposition 17.7 coefficient (Part IV), an iteration potential with optimal leading cost (Part V). 07: transport-sensitive rigidity at Lemma 10.9, with orders five and nine and a sharp threshold (Part II). 08: the fourth-order cube theory with the complete odd-order three-cube polynomial and the refutation of the quartic lower bound (Part III). 09: the complete classification through order five on every finite abelian group, the sixth-order and determinant-five phenomena (Part III). 10: the optimal density cap and the frontier (Part I), a weighted energy extraction and the vertical-fibre Freiman extraction with $2^{-209}\gamma^{103}$ (Part II), torsion-sensitive counting and the polynomial-colouring threshold (Part III), iteration bookkeeping (Part V).

Overlaps. The ten are independent manuscripts: no two of them share more than 1.14% of their text (intake measurement, 8-word shingles). Sources 06 (*Gowers_Refinements (1).zip*) and 10 (*gowers_refinements.zip*) are not editions of one another (1.09%; different titles and theorem sets).

Where the second write had to choose. Each delivered section of 06–10 is printed whole in exactly one Part, by its main subject; a section that serves several themes goes where its main result is used: source 06’s computation section (Section 7) is in Part IV with its certified carry counts, source 09’s Section 13 (relation to progression bounds) in Part III with its cube results, sources 08’s and 09’s certificate appendices beside the theorems they certify, their provenance appendices in Part V. Source 07 is a chapter of Part II with its notation table first, by the intake’s recommendation; source 09 is printed after source 08 in Part III as the fullest cube treatment, and sources 03 and 04 are not moved. Results proved by several sources are printed in each source’s form with dated credit notes at both places, as in the first write.

[write] **6 October 2026; sources 11–14.** *Sources 11–14.* Three archives arrived in commit 08ab4187e and one in d179062cc (both 6 October 2026); 6f530c29c (“Place batch 116”) placed them in this directory as prefixed files and retired them from docs/incoming/. They survive in the arrival commits (`git show 08ab4187e:docs/incoming/<archive>`, `git show d179062cc:docs/incoming/<archive>`). Not staged: the four PDFs, the four manuscripts and delivery READMEs, whose text is printed here, and two checksum manifests (12: 11/11, 14: 13/13 verified at intake). The intake record is the dossier `dossier116_GOWERS`.

No.	Archive (arrival)	Bytes	Files	Pin	Manuscript
11	Gowers_Affine_Rigidity_Research.zip (08ab4187e)	406,963	10	40f885098	1,259 lines, 19 pp.
12	gowers_bohr_extension.zip (08ab4187e)	550,113	12	9ea383e9b	1,215 lines, 25 pp.
13	gowers_local_refinements.zip (08ab4187e)	831,549	17	9ea383e9b	1,223 lines, 31 pp.
14	gowers_extremal_refinements.zip (d179062cc)	518,912	14	9ea383e9b	2,767 lines, 37 pp.

Source 13 reuses the archive name of source 03. The file `docs/incoming/gowers_local_refinements.zip` was added in `66f24b0d0` (source 03, 408,961 bytes), removed in `aefb0b44a` and added again, with different content, in `08ab4187e` (source 13, 831,549 bytes); each is retrieved from its own arrival commit. The two manuscripts are independent (different titles, pins and files; 0.97 % shared text, intake measurement). No two of the fourteen sources share more than 1.14 % of their text, and every overlap involving sources 11–14 is below 1 %.

Pins. `40f885098` (source 11) and `9ea383e9b` (sources 12–14) are commits of 6 October 2026 and ancestors of HEAD; every repository blob they cite was checked at intake and is unchanged at HEAD, except the ledger and `FORMALIZATION_STATUS.txt`, which have moved (Section F.5).

Authorship and disclosures. Source 11: “Research article prepared for Vladimir Reshetnikov” (title block and PDF author field; running head “Research note”). Source 12: “Research article and reproducibility supplement”; it says the repository was “inspected through the GitHub connection”. Source 13: “Prepared for Vladimir Reshetnikov” on its title page, PDF author field “Research report prepared for Vladimir Reshetnikov”; its snapshot was inspected “through the user-selected GitHub connection”. *Source 14 states ChatGPT authorship:* its title page reads “Research report prepared by ChatGPT” and its PDF author field “ChatGPT, prepared for Vladimir Reshetnikov”; its Appendix B mentions an internal review “within the preparation process”, which is not part of its delivery. No source names a human author. Source 14 calls sources 04 and 10 of this report “earlier unpublished project reports” (its references [58] and [59]).

What each contributed. 11: affine rigidity near the endpoint of Lemma 6.2, with exact cubic profiles, a correction step, the exact largest non-affine graph energy and a high-correlation corollary of Proposition 6.1 (Part II). 12: sharp Freiman extension, sharp majority correction, finite Bohr-radius selection and fourth-moment smoothing, giving a polynomial radius in Theorem 10.13 (Part II). 13: prescribed-length simultaneous partitions and obstructions, selection and joint frontiers, a Fourier increment (Part I), the centered three-cube orbit moments, interval discrepancy and telescoping (Part III). 14: the exact secondary cube coefficient and the exponent-five phenomenon, exact polynomial-colouring supersaturation (Part III), coset stability for near-maximal energy and cube collisions and arrangement equality (Part II).

Where the third write had to choose. Each delivered section is printed whole in exactly one Part, by its main subject; conventions sections open the source’s first chapter that uses them (13’s Section 2 in Part I); certificate appendices stand beside what they certify (12’s Appendix A in Part II, 14’s Appendix A in Part III), reproducibility appendices in Part V. Source 14’s examples section (its Section 7) is in Part III with its cube and colouring results, although one subsection illustrates its Part II stability theorem. Source 13’s results that duplicate earlier sources (its Sections 4–6 and 9 in the main) are printed in full with dated credit notes at both ends, as in the earlier writes; no proof was replaced by a pointer.

[write] Independent check of the write for sources 01–05, 6 October 2026. An adversarial check made by the intake after the first write (commit `18507e2b2`) rebuilt that text (167 pages, byte-identical to its PDF), re-derived the four statements proved in the write, re-checked all 56 file and line citations of the crosswalk and notes (all correct, and unchanged at `9f83dbe0f`), the ledger counts, the corrected-statement comparisons, the merge fidelity (all 67 delivered sections present and token-identical apart from the declared interventions; all 425 delivered labels present) and the question counts, and found no refutation of any write claim and no unproved claim. It found four wording defects, corrected at the third write with the first wording kept in dated notes: the comparison of the three Fourier increments said that source 01’s

Theorem I.6.1 is “the strongest”, which is false for lengths at finite N (source 05’s exact length can exceed it by one; example $N = 50$), and called the trade-off lengths “much shorter” where the loss is a constant factor (note after Table 3); the crosswalk row “Lemma 3.8” listed the Minkowski tool of Lemma 3.9 and omitted `lemma_3_9_holds`; and the stronger pair of `cor25_large_scale` exists only inside its proof. Three optional precisions were also applied (status-table sentence, the Corollary 2.4 row, and two wordings), and the first write’s README count of delivered labels (424) is corrected to 425 in the README.

[write] **Independent check of the write for sources 06–10, 6 October 2026.** An adversarial check made by the intake after the second write (commit `9f83dbe0f`) rebuilt that text (321 pages, no warnings), recomputed Remark W1’s counterexample, re-checked source 07’s order five and order nine against `lemma_10_9` and `lemma_10_6`, the five statements proved in that write, the merge fidelity (all sections of sources 06–10 present; all 408 delivered labels; the 50 questions and their credits) and the certificates ($C_4 = 154$ twice; source 09’s determinant-five configuration; $2^{-209}\gamma^{103}$), and found no mathematical error. It found five wording defects, corrected at the third write with the first wording kept in dated notes: the one-paragraph summary of source 07’s chapter named only one of the order-nine requirements (Section II.11); the summary in Section V.2 said that source 01’s increment dominates source 10’s by $3/2$ “only for $\gamma \leq 1/2$ ”, which is false as worded (the bound holds exactly when $\gamma(4 - (2 - \delta)\kappa) \leq 2$, which extends past $1/2$); the progression-threshold note cited source 06’s homogeneous corollary for a criterion that is its existence corollary; the implication from source 10’s Lemma 7.5 replacement used $C \geq 1$, which needs $B \neq \emptyset$ (the empty case is trivial); and the summary of answered questions over-read the finite-cell item. (The check’s further point, that this summary omits source 04’s Question 10, does not apply: the summary records it.)

[write] **6 October 2026; sources 15–20.** *Sources 15–20.* The archives arrived in five commits of 6 October 2026: `212999d4b` (15, 14:22), `82bc38a30` (18 and 19, 14:45), `99f28708e` (16, 14:48), `060528502` (17, 14:55) and `83f3171c1` (20, 15:03). `eed2ab863` (“Place batch 117”) placed 15–17 and `58aa0c493` (“Place batch 118”) placed 18–20 in this directory as prefixed files and retired the archives from `docs/incoming/`; they survive in the arrival commits (`git show <arrival>:docs/incoming/<archive>`). The numbering follows the two batches (each in arrival order, then ASCII order of archive name), so 18 and 19 arrived between 15 and 16; the placement commits record this choice. Not staged: the six PDFs, the six manuscripts and delivery READMEs, whose text is printed here, and four checksum manifests (16: 14/14, 18: 9/9, 19: 12/12, 20: 9/9 verified at intake). The intake records are the dossiers `dossier117_GOWERS` and `dossier118_GOWERS`.

No.	Archive (arrival)	Bytes	Files	Pin	Manuscript
15	<code>collision_sensitive_arrangements.zip</code> (212999d4b)	487,136	9	<code>9ea383e9b</code>	1,752 lines, 27 pp.
16	<code>Gowers_Localization_and_Quantitative_Refinements.zip</code> (99f28708e)	669,627	15	<code>dca9f0638</code>	2,127 lines, 30 pp.
17	<code>Gowers_Polynomial_Restriction_Research.zip</code> (060528502)	427,352	7	<code>dca9f0638</code>	1,327 lines, 21 pp.
18	<code>gowers_cubes_degeneracy_partitions.zip</code> (82bc38a30)	685,030	10	<code>dca9f0638</code>	2,672 lines, 35 pp.
19	<code>gowers_interpolation_article.zip</code> (82bc38a30)	752,841	13	<code>dca9f0638</code>	2,505 lines, 32 pp.
20	<code>gowers_quadratic_counting.zip</code> (83f3171c1)	483,898	10	<code>(1bd57e91c)</code>	1,535 lines, 21 pp.

Pins and what they could see. `9ea383e9b` (13:35, source 15) predates every placement in this directory, so source 15 saw neither source 04 nor source 14. `dca9f0638` (14:17, sources 16–19) and `1bd57e91c` (14:31, source 20’s observed checkpoint) contain the placements of sources 01–10, with source 04’s manuscript staged as `article.tex`, but not the first write (14:34): those sources could read source 04, and none of sources 01–03 and 05–14. All three are ancestors of HEAD; every repository blob they cite was checked at intake and is unchanged at HEAD except the ledger, `FORMALIZATION_STATUS.txt`, `Section16.lean` (source 19) and this report’s `article.tex` (source 20). No two of the twenty sources share more than 1.14% of their text; every

overlap involving sources 15–20 is at most 1.06 % (intake, 8-word shingles), and sources 18 and 19, written against the same pin, prove the same partition theorem in different words.

Authorship and disclosures. Source 15: “Research manuscript prepared for Vladimir Reshetnikov” (title block and PDF author field). Source 16: “Research report prepared for Vladimir Reshetnikov” (both). Source 17: “Research note prepared for the ProveIt project” (both). Source 18: “Research report prepared for Vladimir Reshetnikov” (both); its Appendix A says that “the core derivations were cross-reviewed independently within this task”. Source 19: “Prepared for Vladimir Reshetnikov” on its title page, while its PDF author field reads “ChatGPT, research report prepared for Vladimir Reshetnikov”; its source audit records the retrieval method “GitHub plugin read tools”, and its Section V.81 mentions a written-proof review independent of the initial derivation. Source 20: “Prepared for Vladimir Reshetnikov” (both). Sources 15, 16, 17, 18 and 20 contain no statement about AI assistance; none is invented here. No source names a human author. The reviews mentioned by sources 18 and 19 are not part of their deliveries, and nothing here rests on them.

What each contributed. 15: collision-sensitive amplification coupled with the product property at Lemmas 14.6–14.8, with Cartesian stability and equality results (Part II). 16: Proposition 17.7 without carry patterns (Part IV), own-step partitions for affine frequencies (Part I), and cube coefficients (Part III). 17: polynomial-loss restriction at Section 9, density-sensitive moments and fixed-characteristic systems (Part II). 18: the sharp degenerate-arrangement density at Lemmas 15.4–15.5 (Part II), the cube expansion through order five (Part III) and polynomial partitions by descent in degree with an obstruction (Part I). 19: common sampling, an exact minimax and shared-slice interpolation at Lemma 16.10 (Part II), and degree-profile partitions with an obstruction (Part I). 20: two-function counting, sharp torsion constants and quadratic counting errors with indicator sharpness (Part III).

Where the fourth write had to choose. Each delivered section is printed whole in exactly one Part, by its main subject. Source 16’s conventions section opens its Part IV chapter, the first that uses it. Source 18’s Section 6 (quantitative consequences) is in Part V, like source 04’s Section 9 and source 10’s Section 10; source 19’s Section 8 (worked comparisons) stands with its Section 16 chapter in Part II. The intake recommended printing several results of source 15 that duplicate sources 04 and 14 as pointers; the report keeps its convention instead: every section is printed whole, and the duplicates carry dated credit notes at both ends saying whether the proof is the same argument and whether it is weaker (source 15’s stability and spectral theorems are weaker second routes).

[write] 6 October 2026; sources 21–27. *Sources 21–27.* The archives arrived in five commits of 6 October 2026: 26908b3b6 (21, 22 and 23, 15:27), e2bf630b0 (24, 15:32), d7c4aa399 (25, 15:37), 7e827c6f3 (26, 15:51) and 6078f6e16 (27, 16:16). The merge f7efeb4f5 dropped the three archives of 26908b3b6 from its tree and e3125e698 restored byte-identical copies (the placement records this). a902613c9 (“Place batch 119”) placed 21–23 and 79d9075b9 (“Place batch 120”) placed 24–27 in this directory as prefixed files and retired the archives from docs/incoming/; they survive in the arrival commits (git show <arrival>:docs/incoming/<archive>; for 21–23 also e3125e698). Sources 21–23 arrived in one commit and are numbered in ASCII order of archive name; 24–27 are numbered by arrival. Archive 27 is gowers_arrangement_degeneracy(2).zip: the suffix is a download name, and 27 is an independent manuscript, not an edition of 25 (0.57 % shared text, different title, sources and results). Not staged: the seven PDFs, the seven manuscripts and delivery READMEs, whose text is printed here, three checksum manifests (22: 13/13, 23: 16/16, 26: 8/8 verified at intake) and source 23’s data/latex_build.log (a build log). The intake records are the dossiers dossier119_GOWERS and dossier120_GOWERS.

No.	Archive (arrival)	Bytes	Files	Pin	Manuscript
21	<code>gowers_characteristic_five.zip</code> (26908b3b6)	815,428	19	43306653a	1,775 lines, 24 pp.
22	<code>gowers_degree_stratification.zip</code> (26908b3b6)	391,101	14	43306653a	2,071 lines, 30 pp.
23	<code>gowers_sharp_thresholds.zip</code> (26908b3b6)	832,508	17	65f4e039f	2,320 lines, 33 pp.
24	<code>gowers_section7_refinements.zip</code> (e2bf630b0)	484,922	7	2e0553b0c	1,711 lines, 22 pp.
25	<code>gowers_arrangement_degeneracy.zip</code> (d7c4aa399)	444,249	10	568d1802e	1,207 lines, 17 pp.
26	<code>Report275_Threshold_Free_Affine_Density_Increments.zip</code> (7e827c6f3)	405,485	9	c94998a5d	494 lines, 12 pp.
27	<code>gowers_arrangement_degeneracy(2).zip</code> (6078f6e16)	518,856	9	none	1,824 lines, 28 pp.

Pins and what they could see. 2e0553b0c (13:44, source 24) predates every placement in this directory: source 24 could see neither source 03 nor source 10, whose results it largely repeats. 43306653a (14:45, sources 21 and 22), 65f4e039f (14:48, source 23), 568d1802e (14:52, source 25) and c94998a5d (15:01, source 26) contain the first write (sources 01–05) as `article.tex`; source 18’s archive arrived at 14:45 and sat unopened, and the manuscripts of sources 13, 14 and 18 reached sources 21–23 as separate files (their bibliographies say so), not through the repository. Source 27 records blob identifiers only; it cites neither source 18 nor source 22, whose texts were not in the repository. All pins are ancestors of HEAD; every Lean file the seven sources cite (`Sections01_03`, `Proofs03CubeUpper`, `Section05`, `Definitions`, `Proofs05Corollary58`, `Sections06_07`, `Proofs07DRC`, `Proofs07BalogSzemerédi`, `Sections14_15`) and the edited paper have the same blob at HEAD 2dfe46f44 as at the pins; the ledger, `FORMALIZATION_STATUS.txt` and this report’s `article.tex` have changed. No two of sources 21–27 share more than 1.53% of their text (21 and 23), and none shares more than 2.58% with an earlier source (21 and 14; intake, 8-word shingles); source 26 shares 4.70% with its later sibling, source 32 (template sentences and pinned links). Sources 21 and 23 prove the same theorem in different words, from pins three minutes apart, and neither cites the other; sources 25 and 27, and source 22 at $d = 1$, likewise re-derive source 18’s results independently.

Authorship and disclosures. *Source 21:* title page “Mathematical research report prepared by ChatGPT”, PDF author field “ChatGPT; prepared for Vladimir Reshetnikov”; its `data/mathematical_review.txt` (shipped) records an “AI-assisted” review, “not external peer review”. *Source 22:* “Research report developed with ChatGPT”, PDF author field “ChatGPT; prepared for Vladimir Reshetnikov”. *Source 23:* “Research article prepared for integration into the ProveIt repository” (PDF author field “Research note prepared for ProveIt”), with no statement about AI assistance. *Source 24:* “Research draft prepared with ChatGPT” (both). *Source 25:* “Mathematical development and computational checks with ChatGPT” (PDF author field “Research prepared with ChatGPT for Vladimir Reshetnikov”). *Source 26:* “Research report prepared for Vladimir” (both), with no AI statement. *Source 27:* “An AI-assisted research draft” (title page; PDF author field “AI-assisted research draft prepared for Vladimir Reshetnikov”), naming no system. No statement is invented where a source makes none, and no source names a human author. Source 21 mentions “an independent mathematical audit” and source 22 an “independent internal review”; nothing here rests on them.

What each contributed. 21: the sharp exponent-five coefficient with a uniform remainder, its phase envelope and rigidity, the next coefficient on $\mathbb{Z}/5\mathbb{Z}$ and indicator sharpness (Part III). 22: a degree-stratified count of degenerate arrangements and the case $d = 1$ (Part II); exact pattern minima and obstructions for polynomial colourings (Part III). 23: sharp prescribed-length partition thresholds (Part I), a second proof of 21’s theorem (Part III), and a common base partition for Lemma 16.10 (Part II). 24: a second treatment of Section 7 with an extraction for $\mathbb{F}_p^d$ codomains (Part II). 25: the codimension-one geometry of degenerate arrangements and a selection theorem for every d (Part II). 26: Corollary 5.8 at degree one without a scale threshold (Part I). 27: the second-order structure of degenerate arrangements and an exact count (Part II).

Where the fifth write had to choose. Each delivered section is printed whole in exactly one Part, by its main subject. Sources 21 and 23 prove the same theorem: source 21 is printed as the chapter (its constants are sharper, $42\sqrt{3}$ against 84, and it has the phase envelope and the proof that $20/9$ is the next coefficient on $\mathbb{Z}/5\mathbb{Z}$), and source 23’s Section 7 and its Appendix A, which computes $20/9$ for one model only, follow inside that chapter as a credited second route. Source 23’s three subjects go to three Parts; its Section 8 joins the Section 16 material of Part II. Source 24’s sections are printed whole although most of them repeat sources 03 and 10 with equal or weaker constants (dated notes at both ends say so). Source 26’s Appendix A (integer boxes) stays with its Part I chapter, with a pointer from the Section 16 guide. Sources 25 and 27 present the leading degeneracy coefficient as their own; source 18 is credited as its first print at both ends, and their genuinely new results are named in the guide.

[write] **6 October 2026; sources 28–33.** *Sources 28–33.* The archives arrived in four commits of 6 October 2026: 624d7a42a (28 and 29, 16:05), 12da35673 (30 and 31, 16:07), 6c2e2172b (32, 16:11) and 17dd7886c (33, 16:13). c5513046e (“Place batch 121”) placed them in this directory as prefixed files and retired the archives from docs/incoming/; they survive in the arrival commits (`git show <arrival>:docs/incoming/<archive>`). They are numbered by arrival and then in ASCII order of archive name. Not staged: the six PDFs, the six manuscripts and delivery READMEs, whose text is printed here, four checksum manifests (29: 18/18, 31: 13/13, 32: 8/8, 33’s `manifest.json`: 18/18, verified at intake) and source 33’s two `.log` files (the stdout of its verifier and a byte copy of its erasure JSON). The intake record is the dossier `dossier121_GOWERS`.

No.	Archive (arrival)	Bytes	Files	Pin	Manuscript
28	<code>gowers_logarithmic_selection.zip</code> (624d7a42a)	650,769	14	e3125e698	2,594 lines, 36 pp.
29	<code>gowers_moments_stability_restriction.zip</code> (624d7a42a)	725,320	19	e3125e698	2,256 lines, 30 pp.
30	<code>Gowers_Fixed_Radius_Bohr_Invariance.zip</code> (12da35673)	414,517	10	8d65d3c9f	1,278 lines, 19 pp.
31	<code>gowers_weighted_spectrum_rigidity.zip</code> (12da35673)	371,304	14	e3125e698	1,745 lines, 25 pp.
32	<code>Report276_Threshold_Free_Quadratic_Density_Increments.zip</code> (6c2e2172b)	410,590	9	c94998a5d	533 lines, 12 pp.
33	<code>Gowers_Respected_Arrangement_Stability.zip</code> (17dd7886c)	481,838	19	9f83dbe0f	1,623 lines, 24 pp.

Pins and what they could see. c94998a5d (15:01, source 32, the pin of its sibling source 26), e3125e698 (15:28, sources 28, 29 and 31) and 8d65d3c9f (15:41, source 30) contain the first write (sources 01–05) as `article.tex`; 9f83dbe0f (15:37, source 33) is the second write (sources 01–10). Sources 28 and 29 read the manuscripts of later sources from their archives (28: 07 and 12; 29: the eighteen archives of 06–23, recorded in its `data/source_audit.json`); source 31 read the incoming sources 21–23; source 30 read the README from the moving branch (its `SOURCE_AUDIT.md` says so). All pins are ancestors of HEAD; every Lean file the six sources cite (`Definitions`, `Section05`, `Section10`, `Sections06_07`, `Proofs07FreimanRestriction`) and the edited paper have the same blob at HEAD 5d4718772 as at the pins; the ledger, `FORMALIZATION_STATUS.txt` and this report’s `article.tex` have changed. Source 32 shares 4.70 % of its text with its sibling, source 26 (template sentences and the shared variance section); source 33 shares 1.11 % with source 11; every other pair of sources 28–33 and every pair with an earlier source shares at most 0.93 % (intake, 8-word shingles). Three pairs are independent same-day overlaps: sources 28 and 29 (bounds for c_d^{odd}), sources 29 and 31 (the Proposition 6.1 endpoint), all from one pin, neither citing the other; and sources 30 and 31 repeat parts of sources 12 and 10, which they do not cite.

Authorship and disclosures. *Source 28:* title page “Mathematical development and computational checks with ChatGPT”, PDF author field “Research report prepared with ChatGPT for Vladimir Reshetnikov”; its Section V.112 says the manuscript was “separately checked by other reasoning agents” and not peer

reviewed. *Source 29*: “Mathematical development and verification with ChatGPT”, PDF author field “ChatGPT; prepared for Vladimir Reshetnikov”. *Source 30*: “Mathematical development and exact checks with ChatGPT”, PDF author field “Research report prepared with ChatGPT”. *Source 31*: “Mathematical development and verification with ChatGPT”, PDF author field “Research developed with ChatGPT for Vladimir Reshetnikov”; its `VALIDATION.json` (shipped as `data/31-weighted-spectra-VALIDATION.json`) records “internal AI-assisted reviews” and no independent human peer review. *Source 32*: “Research report prepared for Vladimir” (both), with no AI statement. *Source 33*: “Mathematical development and exact computational checks with ChatGPT”, PDF author field “Research report prepared with ChatGPT for Vladimir Reshetnikov”; its title page says that no human authorship or independent referee review is implied. No statement is invented where a source makes none; nothing here rests on the reviews the sources mention.

What each contributed. 28: the logarithmic selection at Section 10 (Part II); the lower witness for c_4^{odd} and the failure of single-cosine extremality (Part III). 29: central-binomial moment bounds for c_d^{odd} (Part III); the joint endpoint with soft and mixed selectors and the fixed-field restriction (Part II). 30: the fixed-radius escape inequality and the exponential-rank obstruction (Part II). 31: the sharp near-extremal U^3 profile and graph stability; weights through the Section 7 restriction (Part II). 32: Corollary 5.8 at degree two without a scale threshold (Part I). 33: the linear branch at Lemma 12.3 and sharp high-agreement repair (Part II).

Where the sixth write had to choose. Each delivered section is printed whole in exactly one Part, by its main subject. For c_d^{odd} source 29 is the main source (its upper bound supersedes source 28’s at every $d \geq 4$), and source 28’s Section 8 follows it with its lower bound and obstruction. At the Proposition 6.1 endpoint source 31 is the main source (its constants are sharp), and source 29’s Sections 5–8 follow as a second route that also has soft and mixed selectors. At Section 7 source 10 stays the spine; source 31’s Sections 6–8 are printed whole though they repeat source 10 with weaker or equal constants, with credits at both ends. At Section 10 sources 07 and 12 stay the spines for paths, extension and synchronization; source 30’s extension and compatibility steps are credited to source 12. Source 32 joins source 26’s chapter, as planned there; its Sections 6–7 go to Part V like source 26’s.

[write] **6 October 2026; sources 34–39.** *Sources 34–39.* The archives arrived in three commits of 6 October 2026: 70c6a1cf1 (34, 35 and 36, 17:03), a33bcaedf (37, 17:06) and b103a65d4 (38 and 39, 17:11). bf2fe146d (“Place batch 122”) placed them in this directory as prefixed files and retired the archives from `docs/incoming/`; they survive in the arrival commits (`git show <arrival>:docs/incoming/<archive>`). They are numbered by arrival and then in ASCII order of archive name. Not staged: the six PDFs, the six manuscripts and delivery READMEs, whose text is printed here, two checksum manifests (34: 16/16, 38: 32/32, verified at intake) and source 38’s `data/latex_build.log` (the TeX log of its sandbox build). The intake record is the dossier `dossier122_GOWERS`.

No.	Archive (arrival)	Bytes	Files	Pin	Manuscript
34	Gowers_Cube_Torsion_and_Norm_Gap s.zip (70c6a1cf1)	607,948	17	12da35673	2,805 lines, 37 pp.
35	Gowers_Graded_Refinements_TeX_an d_PDF.zip (70c6a1cf1)	662,283	13	12da35673	3,007 lines, 40 pp.
36	Gowers_Sharp_Symmetry_Certificat ion.zip (70c6a1cf1)	487,472	10	6c2e2172b	1,555 lines, 19 pp.
37	Gowers_Sharp_Missing_Intervals.z ip (a33bcaedf)	457,872	10	869289aa3	1,616 lines, 27 pp.
38	gowers_quotients_rigidity.zip (b103a65d4)	681,853	33	12da35673	3,217 lines, 46 pp.
39	ramsey_shared_anchor_interpolati on.zip (b103a65d4)	425,241	15	869289aa3	2,164 lines, 31 pp.

Pins and what they could see. 12da35673 (16:07, sources 34, 35 and 38), 6c2e2172b (16:11, source 36) and 869289aa3 (16:17, sources 37 and 39) contain the second write (sources 01–10) as `article.tex` (blob 3adb272), with sources 11–33 placed or in `docs/incoming/`. Source 34 read the incoming archives of sources 29 and 31, not source 28’s; source 35 read source 40’s archive and those of sources 28 and 29, and the

placed sources 18, 22 and 23; source 38 read sources 21, 22, 23 and 31 (archive and manuscript hashes in its `SOURCE_LEDGER.json`); source 36 read the transcription, `FORMALIZATION_STATUS.txt` and the README and says that the large merged article and the pending reports were not exhaustively read; sources 37 and 39 read Lean files of Sections 5 and 16, and the README from the moving branch (both say so; the blob is the pinned one), and source 39 no research report. All pins are ancestors of HEAD; every Lean and paper file the six sources cite has the same blob at HEAD 89c1cbc38 as at the pins (56 provenance records, all matching, rerun at the write); the ledger, `FORMALIZATION_STATUS.txt`, the README and this report’s `article.tex` have changed. The largest textual overlaps are source 38 with source 21 (2.60 %) and with source 31 (2.11 %), where source 38 reproves their background in its own words, then source 35 with source 28 (1.00 %) and with source 40 (0.91 %); every other pair involving sources 34–39 shares at most 0.78 %, and source 39 shares less than 0.5 % with source 19, whose results it re-derives (intake, 8-word shingles). Same-day overlaps without citation: sources 34 and 35 (the U^4 coefficient, from one pin); source 34 with source 28 (the two-harmonic example and dual identity); source 36 with sources 31 and 29; source 39 with source 19.

Authorship and disclosures. *Source 34*: title page “Mathematical development and exact computation with ChatGPT”, PDF author field “ChatGPT; research report prepared for Vladimir Reshetnikov”; its Section 1 says “with internal proof review”; its delivery README has no AI statement. *Source 35*: “Mathematical development and computational checks with ChatGPT”, PDF author field “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov”; its README calls the manuscript AI-assisted and unrefereed. *Source 36*: “Research manuscript prepared with ChatGPT” (both); its README calls it an unrefereed AI-assisted research draft, not independently reviewed, and its `SOURCE_AUDIT.md` calls its checks self-checks. *Source 37*: “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov” (both). *Source 38*: “Mathematical development and verification with ChatGPT”, PDF author field “Mathematical research developed with ChatGPT for Vladimir Reshetnikov”; its Section 1 says that its independent checks are separate AI derivations and computations. *Source 39*: “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov” (both). No statement is invented where a source makes none; nothing here rests on the reviews the sources mention.

What each contributed. 34: the prime-torsion onset and its asymptotics, positivity, consecutive binary determinants, $c_4^{\text{odd}} \leq 2/3$ (Part III); the sharp nonuniform transport threshold (Part II). 35: degree-weighted arrangement selection (Part II); the coefficient 5 and a U^3 energy inequality (Part III); polynomial divisibility and prescribed-length thresholds (Part I). 36: sharp symmetry certification and lossless polarization (Part IV). 37: the missing-interval minimax, Lemma 5.2 and a recurrence exponent (Part I). 38: exact quotient reduction in exponent five (Part III), the sharp odd-order stability modulus (Part II), six-term defects (Part III). 39: shared-anchor sampling and labelled refinement at Lemma 16.10 (Part II).

Where the seventh write had to choose. Each delivered section is printed whole in exactly one Part, by its main subject. For c_4^{odd} source 34 is the main source (its coefficient supersedes source 35’s), source 35’s Section 6 follows it, and both are appended to the existing chapter rather than given a new one. For the characteristic-five profile source 38 is the main source for the dimension-uniform statement, source 21 for the first coefficients; for the Proposition 6.1 endpoint source 31 stays the spine and source 38 adds odd order. For Lemma 16.10 source 19 stays the spine; source 39 is printed whole although much of it re-derives source 19, with credits at both ends. Source 35’s Part 1 builds on source 40, which is placed and not yet written: the forward reference is kept and dated. Source 34’s Appendix A is printed as its own short chapter beside source 07’s, not inside it, because no number printed earlier may move.

[write] 7 October 2026; sources 40–46. *Sources 40–46.* The archives arrived in six commits of 6 October 2026: 13f105f42 (40, 15:31), a6eea2a19 (41 and 42, 17:40), 7382d8c44 (43, 17:49), 0b4890a1c (44, 17:50), f70d42385 (45, 17:55) and f2baec573 (46, 17:56). 2dfe46f44 (“Place batch 123”) placed them in this directory as prefixed files and retired the archives from `docs/incoming/`; they survive in the arrival commits. Source 40 arrived before batch 120’s first archive and was overlooked when that batch was numbered; it took the next free number at placement, so source 35’s citation of it is a forward reference. Sources 41–46 are numbered by arrival and then in ASCII order of archive name. Not staged: the seven PDFs, the manuscripts (with source 45’s nine section files and `references.bib`) and delivery READMEs, whose text is printed here, and four checksum manifests (41 and 42: 8/8 each, 45: 30/30, 46: 8/8, verified at intake). The intake record is the dossier `dossier123_GOWERS`.

No.	Archive (arrival)	Bytes	Files	Pin	Manuscript
40	fejer-arrangement-selection.zip (13f105f42)	427,478	7	none	1,281 lines, 20 pp.
41	Report277_Threshold_Free_All_Deg ree_Density_Increments.zip (a6eea2a19)	527,860	9	c94998a5d	1,022 lines, 21 pp.
42	Report278_Sharp_Quadratic_Scale_ Density_Increments.zip (a6eea2a19)	421,797	9	b16ce380d	688 lines, 15 pp.
43	Boolean_Phase_Integration_Resear ch.zip (7382d8c44)	523,740	9	797aa3cfe	1,938 lines, 26 pp.
44	gowers_grid_core_research.zip (0b4890a1c)	513,096	13	8c12c906e	1,632 lines, 24 pp.
45	Gowers_Rigidity_Gaps_Research.zi p (f70d42385)	695,194	31	6ce33539e	3,589 lines in 10 files, 45 pp.
46	gowers_polarization_barriers.zip (f2baec573)	498,767	9	none	1,503 lines, 23 pp.

Pins and what they could see. Source 40 records no pin and no blob: it read the transcription’s Section 15 and the README of this directory, by path. c94998a5d (15:01, source 41) has `Section05.lean` and `Definitions.lean` with the blobs 882b17a and 97113b7 that source 41 records; it also records, to correct the staged `SOURCES.md` of sources 26 and 32, the raw SHA-256 of four files whose recorded values hash the file plus one LF (verified). b16ce380d (16:40, source 42) contains this report at its third write (sources 01–14), with source 10’s variance-to-density calculation at the lines source 42 cites; source 42 identifies sources 41 and 26 by the SHA-256 of their manuscripts (both match the delivered files). 797aa3cfe (17:06, source 43): `Sections17_18.lean` and `Definitions.lean`, and this directory’s README as observed on the main branch; it read source 36 from the user’s library (by its question labels). 8c12c906e (17:11, source 44): `Sections14_15.lean` and the transcription; sources 25 and 27 from the user’s library (“Baseline A” and “Baseline B”). 6ce33539e (17:05, source 45): eight blobs, among them this report’s `article.tex` at its third write and source 15’s integration note and verifier. Source 46 records the blobs of the transcription and of this directory’s README (41085a0, fe4d31eb), no commit. All pins are ancestors of HEAD; every Lean and paper file the seven sources cite has the same blob at HEAD 6d4942222 as at the pins (30 provenance records, all matching, rerun at the write); the ledger, `FORMALIZATION_STATUS.txt`, the README and this report’s `article.tex` have changed. Textual overlaps (intake, 8-word shingles): source 41 contains 47.9% of source 32’s shingles, and its Sections 2 and 5 are source 32’s Sections 2 and 4 verbatim (checked at the write); source 41 with source 26 4.4%; source 42 with sources 32, 26 and 41 2.5, 2.1 and 1.9%; source 43 with source 36 1.4%; source 44 with sources 27 and 25 1.2% each; every other pair involving sources 40–46 shares at most 0.9%. Overlaps of results without citation: source 41 with sources 26 and 32 (its retained small-degree bounds and its two-regime method; neither is cited in its article); source 45 with sources 28 and 34 (the single-pair non-extremality and the dual identity) and with sources 29 and 35 (the moment bound), all four in `docs/incoming/` at its pin; source 46 with source 36 (characteristic two against odd characteristic) and with source 43 (the same main theorem, the same day); source 44 cites sources 25 and 27 only by library identifiers.

Authorship and disclosures. Source 40: “Research note prepared for the ProveIt repository” (title page and PDF author field); no statement on AI assistance. Sources 41 and 42: “Report277” and “Report278”, “Research report prepared for Vladimir”; no statement on AI assistance, as for sources 26 and 32. Source 43: “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt project”. Source 44: “An AI-assisted research manuscript prepared for Vladimir Reshetnikov and the ProveIt repository”; its `SOURCE_AUDIT.md` calls its two baselines AI-assisted drafts. Source 45: “Research report prepared with ChatGPT for Vladimir Reshetnikov”; its status statement calls it an AI-assisted research manuscript, and its six `reviews/*.md` (shipped as `45-rigidity-gaps-review_*.md`) describe internal cross-checks by a second AI agent. Source 46: “Research contribution prepared for the ProveIt project”; no statement on AI assistance. No statement is invented where a source makes none; nothing here rests on the reviews the sources mention.

What each contributed. 40: one-stage Fejér selection and the kernel-class barrier (Part II). 41: Corollary 5.8 in every degree (Part I). 42: the arithmetic and quadratic-scale variance increment and the affine square-root cap (Part I). 43: Boolean phase integration, division-free orthogonality, the fourth-power triangle bound and the quartic value (Part IV). 44: core avoidance and the second Whitney coefficient at Lemma 15.4 (Part II). 45: torsion-free energy levels and the coset range (Part II); the U^4 gap 4795/832 (Part III). 46: exact binary cubic norms, the seven-function repair and trace tensors (Part IV).

Where the eighth write had to choose. Each delivered section is printed whole in exactly one Part, by its main subject. For Corollary 5.8 source 41 is the main source in every degree and sources 26 and 32 keep their sharper small-degree constants; source 41 is printed as a chapter of its own at the end of Part I, after source 35's, with source 42 directly after it, because the chapter of sources 26 and 32 is no longer the last of Part I. Source 41's Sections 2 and 5 repeat source 32's Sections 2 and 4 verbatim; the intake recommended printing such duplicated sections once, and the write prints them again, with a note at each, because source 41's later sections refer to their labels and this report replaces no proof by a pointer. For the binary cubic norm source 46 is the main source (its eight-function statement contains the one-function maximum) and source 43 a credited second route printed after it, with its own further results. For c_4^{odd} source 45 is the main source for the upper bound and stands at the end of Part III, because the chapter on the odd-order constants is no longer the last; the lower end is still the batch-121 certificate. Source 40, on which source 35 builds, is printed after source 35's chapter for the same reason.

[write] **7 October 2026; sources 47–53.** *Sources 47–53.* The archives arrived in two commits of 6 October 2026: b6a8ee966 (47–50, 20:13) and 5d4718772 (51–53, 20:20). f8bc5e2ec (“Place batch 124”) placed them in this directory as prefixed files and retired the archives from docs/incoming/; they survive in the arrival commits. They are numbered by arrival and then in ASCII order of archive name. Not staged: the seven PDFs, the manuscripts and delivery READMEs, whose text is printed here, four checksum manifests (49: 12/12, 51: 27/27, 52: 14/14, 53: 32/32, verified at intake), and the 46 Lean source snapshots shipped by sources 51–53 (copies of repository files at their recorded commits, verified byte for byte at intake and again at the write; retrievable from git history). The intake record is the dossier dossier124_GOWERS.

No.	Archive (arrival)	Bytes	Files	Pin	Manuscript
47	Boolean_Obstruction_Energies_Research.zip (b6a8ee966)	469,142	9	5c9a442d2	1,512 lines, 23 pp.
48	Gowers_Collision_Aware_Derandomization.zip (b6a8ee966)	490,403	21	0b4890a1c	1,763 lines, 27 pp.
49	gowers_rank_dual_refinements.zip (b6a8ee966)	666,835	13	8c12c906e	3,329 lines, 45 pp.
50	ramsey_branch_cover_obstructions.zip (b6a8ee966)	517,479	10	588552b79	1,734 lines, 27 pp.
51	Report279_Section13_Square_Geometry_Obstruction_and_Repair.zip (5d4718772)	575,022	28	six, from aad4789fc	948 lines, 22 pp.
52	Report280_Composite_Modulus_Obstructions_to_Stage137.zip (5d4718772)	459,905	15	5c9a442d2	835 lines, 20 pp.
53	Report281_Explicit_Uniform_Density_Square_Extraction.zip (5d4718772)	542,010	33	bf2fe146d	756 lines, 16 pp.

Pins and what they could see. All are ancestors of HEAD and record 101 exact companions and 19 open statements (none of the seven quotes these counts). 5c9a442d2 (18:04; sources 47 and 52): source 47 records the blobs of Sections17_18.lean and Definitions.lean, and read source 43 from the user's library (boolean_phase_integration.tex); source 52 ships five Lean snapshots. 0b4890a1c (17:50, source 48): Proofs09Restriction.lean, and source 17 from the library (gowers_polynomial_restriction.tex, its pin dca9f0638). 8c12c906e (17:11, source 49): this report's article at its third write and the transcription, and the

manuscripts of sources 22, 34, 35 and 19 from the library, identified by SHA-256 in its source ledger (two of six values hash the file followed by one line feed, as for sources 26 and 32). 588552b79 (17:41, source 50): four Lean blobs (`Section16.lean`, `Definitions.lean`, `Sections14_15.lean`, `Proofs16LocalAffineLift.lean`) and this directory’s README. aad4789fc to 4f8cd1198 (17:18–18:00, source 51) and bf2fe146d (18:11, source 53, observed also at 38ba8bd25 and 4f8cd1198): 18 and 23 Lean snapshots. The provenance records were re-verified at HEAD 06b30ca52 (79 records, all matching). *Changed at HEAD*: `Proofs16LocalAffineLift.lean` (source 50; note in its Section II.234), `Sections12_13.lean`, `Proofs13BilinearExtraction.lean`, `Proofs13CommonStepSelection.lean`, `Proofs13CompleteRowExtraction.lean` and `Proofs13AllDensityExtraction.lean` (sources 51–53; notes there), and the transcription, by 21 lines in Section 13 (source 49’s crosswalk table; note there); every other cited file has its pin blob at HEAD. Textual overlaps (intake, 8-word shingles, share of the newer source): source 47 with source 43 2.8%; source 53 with source 51 4.1%, source 52 with source 51 1.9% (the series boilerplate and the shared lemmas of Section II.195); source 48 with source 17 1.0%; source 49 with sources 34 and 35 0.6 and 0.5%; source 50 at most 0.3% with any source. Overlaps of results without citation: source 49 with source 45 (the c_4^{odd} upper bound; 45 postdates 49’s pin); sources 51, 52 and 53 with one another (the same series, the same day); source 53 with source 51 (the rectification, balanced-parent and cover lemmas); sources 47 and 48 cite sources 43 and 17 only by library file names.

Authorship and disclosures. *Source 47*: “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt project”. *Source 48*: “Research manuscript prepared with ChatGPT for the ProveIt project”. *Source 49*: “Research manuscript prepared for Vladimir Reshetnikov with ChatGPT”; its Section V.187 reports an internal review, not peer review. *Source 50*: “Research manuscript prepared for the ProveIt project”; no statement on AI assistance. *Sources 51–53*: “Report279” to “Report281”, “Research report prepared for Vladimir”; no statement on AI assistance, as for sources 26, 32, 41 and 42. No statement is invented where a source makes none.

What each contributed. 47: exact Boolean obstruction energies and symmetric integration thresholds (Part IV). 48: collision-aware derandomization of the restriction step (Part II). 49: the two-cube degeneracy for $d = 1$ and the acyclic Section 16 replacement (Part II); the U^4/U^2 coefficient 29/4 and the U^5 bound (Part III). 50: branch-cover laws and the counterexample to the catalogue lemma_16_10 (Part II). 51: the grid obstruction and the reselection repair at Section 13 (Part II). 52: composite-modulus counterexamples at Stage 13.7 (Part II). 53: the explicit uniform square-extraction threshold with the actual-density side (Part II).

Where the ninth write had to choose. Each delivered section is printed whole in exactly one Part, by its main subject. Source 49 is split by subject: its degeneracy sections join the Lemma 15.4 material in Part II, its norm section Part III, its acyclic cover theorem the Lemma 16.10 chapter, after source 50, whose counterexample is the reason a replacement is needed. For Section 13 source 53 is the spine (its theorem is the strongest positive statement and contains source 51’s repair up to the exponent), followed by source 51 (the obstruction, with its own form of the repair) and source 52 (Stage 13.7); the shared lemmas of sources 51 and 53 are printed in both, because each source’s proofs cite its own labels and this report replaces no proof by a pointer. Source 47 is printed after the binary chapter of sources 46 and 43 in Part IV, and source 48 after source 45’s energy chapter in Part II, because those chapters are no longer the last of their Parts.

[write] **7 October 2026; sources 54–59.** *Sources 54–59.* The archives arrived in two commits of 6 October 2026: aa1d86f6f (54–58, 21:12) and 9458a55a9 (59, 21:44). 7dc1ee29d (“Place batch 125”) placed them in this directory as prefixed files and retired the archives from `docs/incoming/`; they survive in the arrival commits. They are numbered by arrival and then in ASCII order of archive name. Not staged: the six PDFs, the manuscripts (source 56’s main file with its thirteen section files) and delivery READMEs, whose text is printed here, source 55’s packaging `.gitignore`, and three checksum manifests (54: 10/10, 58: 17/17, 59: 14/14, verified at intake). The intake record is the dossier dossier125_GOWERS.

No.	Archive (arrival)	Bytes	Files	Pin	Manuscript
54	Gowers_Single_Phase_Research.zip (aa1d86f6f)	455,454	11	6a6d7961c	1,387 lines, 20 pp.
55	gowers_endpoint_stability.zip (aa1d86f6f)	348,065	12	128f514af	1,498 lines, 21 pp.
56	gowers_relative_affine_repair.zip (aa1d86f6f)	515,881	23	b6a8ee966	1,593 lines in 14 files, 25 pp.
57	gowers_research_proveit.zip (aa1d86f6f)	596,096	11	128f514af	2,340 lines, 31 pp.
58	gowers_torsion_uniform_restricti on.zip (aa1d86f6f)	850,168	18	5d4718772	1,770 lines, 26 pp.
59	gowers-character-affine-refinem ents.zip (9458a55a9)	582,008	15	bec460a33	1,878 lines, 32 pp.

Pins and what they could see. All are ancestors of HEAD. The ledger recorded 103 exact companions and 17 open statements at 6a6d7961c (19:52, source 54) and 105 and 15 at the other pins (20:13–20:29); `theorem_13_12` was open at all of them, and none of the six quotes these counts. At every pin the article held sources 01–27 (01–33 at bec460a33, source 59’s pin, the sixth write); sources 28–46 were placed but not yet written when sources 54–58 were pinned. 6a6d7961c (source 54): the transcription, `Sections08_09.lean`, source 36’s source audit and the README (blob b2e25fa). 128f514af (sources 55 and 57): source 55 the README and source 45’s energy review, source 57 six Lean blobs (`Definitions.lean`, `Sections12_13.lean`, `Proofs13UniformSquareExtraction.lean`, `Proofs13CompleteSquareExtraction.lean`, `Proofs13ExplicitFourierExponent.lean`, `Proofs13PrintedExponentComparison.lean`), the transcription and `FORMALIZATION_STATUS.txt`, and the archives of sources 40 and 45 at 8ed496b14 (blobs 44036f5, dac681a). b6a8ee966 (source 56): `Sections06_07.lean`, `Section10.lean`, `Proofs10Shift.lean`, the README and the transcription. 5d4718772 (source 58): this article (blob 9897851), the README, `Sections08_09.lean`, `Definitions.lean`, source 29’s audit record and source 40’s audit and verifier (not its manuscript, as it says). bec460a33 (source 59): this article (blob 32087ba), the README and the transcription, by blob and SHA-256. The provenance records were re-verified at HEAD fe7165a3c (pins, every recorded blob, the three checksum manifests and source 57’s recorded SHA-256 of its manuscript and PDF). *Changed at HEAD:* the article and the README (later writes), `FORMALIZATION_STATUS.txt` (source 57 read d846374), the transcription (by 21 lines in Section 13; sources 54, 57 and 59 read blob 41085a0) and `Sections12_13.lean` (source 57 read 1de556b; four Section 13 commits since, 5fff7d9d2 to 8da5b7f07, leave the definitions of `theorem_13_12`, `lemma_13_11`, `secondDifferenceFourier` and `BilinearOn` unchanged, `theorem_13_12` moving from line 626 to line 718); every other cited Lean file has its pin blob at HEAD. Textual overlaps (intake, 8-word shingles, share of the newer source) are at most 1.46 % for every pair (59 with 11 1.46 %, 57 with 11 1.06 %, 54 with 02 0.94 %, 56 with 11 0.85 %, 58 with 17 0.65 %, 55 with 27 0.59 %; with sources 47–53 at most 0.52 %): the overlaps are of results, not text. Overlaps of results without citation: source 55 with sources 31 and 38 (the $d = 3$ endpoint coefficients; both were placed but unwritten at its pin, and the README did not yet describe them); source 54 with source 31’s phase-space identity (its $\mathcal{E}_\eta$ identity is a fibre of 31’s Moyal identity; the same situation); source 58 with source 48 (restriction with repeated entries, the same day; neither cites the other).

Authorship and disclosures. *Source 54:* “Research manuscript prepared for the ProveIt project”; no statement on AI assistance. *Source 55:* “Prepared with ChatGPT for Vladimir Reshetnikov” (title page), “AI-assisted research draft” (status). *Source 56:* “Research draft prepared for the ProveIt project”, “AI-generated research manuscript” (status). *Source 57:* “Prepared for the ProveIt Ramsey project”; its PDF author field “Research draft prepared with ChatGPT for the ProveIt project”, “AI-assisted research manuscript” (baseline statement), and its Section V.217 reports “AI-assisted proof passes”, an internal audit, not peer review. *Source 58:* “Research contribution prepared for the ProveIt project”; its Appendix V.221: “prepared with OpenAI ChatGPT, including independent parallel proof reviews”. *Source 59:* “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov”, “AI-assisted research draft” (status). No statement is invented where a source makes none.

What each contributed. 54: one quadratic phase at Proposition 8.1, offset ceilings and their stability

(Part II). 55: sharp polynomial-phase endpoint stability in every degree (Part II). 56: relative affine repair on Fourier-uniform domains (Part II). 57: Theorem 13.12 in writing, near-maximal frequency extraction, bounded Schur deficits and the third energy level (Part II). 58: Lemma 9.3 on every cyclic modulus and the positive-filter barrier (Part II). 59: exact character separation, target-independent Freiman restriction and the enlarged affine range (Part II).

Where the tenth write had to choose. Each delivered section is printed whole in exactly one Part, by its main subject; the integration, verification, question and conclusion sections and the appendices are in Part V, as before. Sources 57 and 59 are split by subject: 57’s Section 5 joins source 55 at the Proposition 6.1 endpoint, its Sections 2–4 form the Theorem 13.12 chapter and its Section 6 the energy-level chapter; 59’s affine sections (7–10, with its examples and the Green–Tao remark) open the affine chapter, its separation sections (2–6) form the Section 7 chapter. In the affine chapter source 59 is the spine (its local range contains source 11’s and its global corollaries complete source 11’s cutoff), followed by source 56 (the relative branch on sparse uniform domains). The chapters follow the order of Gowers’s sections; source 57’s Section 13 chapter stands after the ninth write’s Section 13 chapter, not inside it, because no number printed earlier may move.

[write] 7 October 2026; sources 60–69. *Sources 60–69.* The archives arrived in two commits: 0d93c1eac (60–64, 6 October 2026, 22:07) and 0281fcb05 (65–69, 7 October 2026, 08:30). 39d3295c4 and 06b30ca52 (“Place batch 126” and “Place batch 127”) placed them in this directory as prefixed files and retired the archives from docs/incoming/; they survive in the arrival commits. They are numbered by arrival and then in ASCII order of archive name. Not staged: the ten PDFs, the manuscripts (with source 67’s front and end matter and four section files, and source 68’s main.tex and eleven section files) and delivery READMEs, whose text is printed here, two run records (61’s checks/run.log, a byte copy of its results file, and 64’s verification.log, which carries a sandbox path), 66’s data/verification.log, and the checksum manifests (60: 18/18, 61: manifest.json 13/13, 63: 29/29, 64: 12/12, 66: 12/12, 67: 31/31, 68: 30/30, verified at intake). The intake records are the dossiers dossier126_GOWERS and dossier127_GOWERS, lost with the session scratchpad of 7 October; every detail used here was re-derived from the archives, extracted afresh from the arrival commits.

No.	Archive (arrival)	Bytes	Files	Pin	Manuscript
60	Gowers_Fourier_Arrangement_Refinements.zip (0d93c1eac)	849,885	19	89c1cbc38	3,145 lines, 49 pp.
61	Gowers_Torsion_Exact_Restriction.zip (0d93c1eac)	355,675	14	5ea1d06dd	1,249 lines, 23 pp.
62	ProveIt_All_Moduli_Restriction.zip (0d93c1eac)	457,236	11	17123532a	1,317 lines, 20 pp.
63	gowers_asymptotic_constants.zip (0d93c1eac)	929,151	30	d7b83fa1f	2,633 lines, 35 pp.
64	gowers_binomial_separation_research.zip (0d93c1eac)	344,942	13	none	1,268 lines, 19 pp.
65	Gowers_All_Degree_Boolean_Integration.zip (0281fcb05)	475,026	10	8c40adc24	1,459 lines, 21 pp.
66	Gowers_Endpoint_Partitions.zip (0281fcb05)	444,292	13	none	1,637 lines, 24 pp.
67	Gowers_Exact_Local_Profiles_TeX_and_PDF.zip (0281fcb05)	870,366	32	none	3,295 lines, 44 pp.
68	Gowers_Sharp_Energy_and_Restriction.zip (0281fcb05)	856,963	31	0d93c1eac	3,422 lines, 44 pp.
69	ramsey_boolean_phase_integration.zip (0281fcb05)	499,619	8	0d93c1eac	1,636 lines, 25 pp.

Pins and what they could see. All pins are ancestors of HEAD; they lie between 21:05 and 22:07 on 6 October 2026, when the ledger recorded 107 exact companions and 13 open statements (theorem_13_12, corollary_5_8 and Lemmas 13.5, 13.6, 13.9 and Corollary 13.10 were still open), and none of the ten quotes

these counts. At 5ea1d06dd, 17123532a, 89c1cbc38 and d7b83fa1f the article held sources 01–33 (the sixth write), at 8c40adc24 and 0d93c1eac sources 01–39 (the seventh); later members were placed but not written. 89c1cbc38 (source 60): the README (blob 689dbd3), this article (32087ba) and the transcription (41085a0). 5ea1d06dd (61): Sections06_07.lean, Sections08_09.lean and Proofs09Restriction.lean by blob. 17123532a (62): the README (689dbd3), Definitions.lean, Sections08_09.lean and Proofs09Restriction.lean. d7b83fa1f (63): the transcription, this article (32087ba) and the README (689dbd3), by blob and SHA-256, and the undelivered predecessor manuscript of its Section 6 by SHA-256 (below). Source 64: Section04.lean (blob 987daf1) and the search snapshots 17123532a, 176e155e5. 8c40adc24 (65): Sections17_18.lean and Definitions.lean, and source 47’s manuscript (its comparison pin 5c9a442d2). Source 66: the transcription (41085a0), the README of the seventh write (1ec6351) and source 45’s energy review (7984f49). Source 67: this article of the seventh write (f4e5705), a README blob cd43de1 that no commit contains: it is the README of the seventh write (1ec6351, 202,687 bytes) with one extra final line feed, which reproduces the recorded blob, size 202,688 and SHA-256 exactly (checked at the write; source 68 records the same artefact for its own copies), the transcription, and three manuscripts “inspected separately” by SHA-256, which are sources 59, 49 and 57 as delivered (verified at the write). 0d93c1eac (68): this article (f4e5705, sources 01–39), the README, the transcription, the ledger, the audit records of sources 58 and 59, and the archives of sources 45, 58 and 59 at 8ed496b14 and f8bc5e2ec. 0d93c1eac (69): Sections17_18.lean (18495f5) and source 47’s manuscript and source audit. The provenance records were re-verified at HEAD fb9602e55 (every pin, every recorded blob, the recorded SHA-256 values of the manuscripts read). *Changed at HEAD among cited files*: the article and the README (later writes), FORMALIZATION_STATUS.txt and the ledger, and the transcription (sources 60, 63, 66, 67 and 68 read blob 41085a0; three Section 13 commits of 7 October, bc87bdf14 to 8da5b7f07, changed it by 23 inserted and 2 deleted lines, and the Lemma 16.10 repair, 456b6e3ad, added four comment lines (an editorial note) before Lemma 16.10; no printed line of the passages these sources cite changed); every cited Lean file has its pin blob at HEAD. Textual overlaps (8-word shingles, share of the newer source, computed at the write) are largest in the Boolean series, where definitions and the reproved canonical calculation recur: 65 with 47 5.76%, 65 with 69 5.06%, 69 with 47 4.90%, 65 and 69 with 43 1.65 and 1.35%; then 68 with 58 2.15% (the reproved expectation theorem) and 60 with 33 1.00%; every other pair is below 0.6%. Apart from those reproved passages the overlaps are of results, not of text. Overlaps of results without citation: source 63 with source 16 (the sliding windows and the geometric ceiling, below); source 64 with sources 10 and 14 (its dimension barrier is source 10’s exact threshold at degree $k - 1$, which source 14 reproves); source 66 with source 57 (the partition normal form and the third level, the same day; 66 read source 45’s review, not 57); source 67 with source 59 at the order-thirteen example (67 credits 59’s ranges, not its $\mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/3\mathbb{Z}$ map); sources 65 and 69 with each other (the same theorem, independently, the same day); source 62 with source 58 (all-moduli restriction, pins before 58’s placement); sources 61 and 59 (the exponent $(2m + 1)\log_2(5m)$, independently).

The predecessor credited by source 63. Source 63’s Section 6 credits “the separate 7 October 2026 manuscript *Local Phase Removal and Cubic Rigidity in Gowers’ Szemerédi Framework*” (126,709 bytes, SHA-256 9cd0a536...) for the sliding-window and positive-cover interface and the previous denominator $\hat{D}_k$. *That manuscript was never delivered*: no file with its SHA-256 is in any of the 966 archive versions ever added under docs/incoming/ or in the HEAD tree (checked at the write), so it is none of the placed members 01–100; source 16, *Local Phase Removal Without Carry Patterns*, is a different manuscript. Source 63 does not cite source 16, whose sliding-window localization and local ceiling $2^{k+1}/(k+2)!$ (Proposition IV.13.5) are those it uses; the credit is recorded in its chapter.

Authorship and disclosures. *Source 60*: “Research prepared with ChatGPT for Vladimir Reshetnikov”. *Source 61*: “Mathematical development and exact computational checks with ChatGPT”, “AI-assisted, unrefereed research manuscript” (status). *Source 62*: “Research manuscript prepared with ChatGPT”, “AI-assisted mathematical research draft” (status). *Source 63*: “Mathematical development and manuscript by ChatGPT”; its development and “separate proof reviews were carried out by ChatGPT instances”, the shipped review notes, “not external peer review”. *Source 64*: “Research manuscript prepared for the ProveIt project”; no statement on AI assistance. *Sources 65, 66, 67 and 69*: “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov” (65, 67 and 69: “and the ProveIt project”, “the ProveIt Ramsey project”). *Source 68*: “Mathematical derivation and manuscript assisted by OpenAI ChatGPT”. No statement

is invented where a source makes none.

What each contributed. 60: the arrangement polynomial and two-direction repair (Part II), the norm limit as a Fourier extremum (Part III), translated restriction on prime moduli (Part II). 61: exact Freiman partitions and threshold-free exact restriction on arbitrary finite abelian groups (Part II). 62: unit-dilation restriction on every cyclic modulus (Part II). 63: the sharp bounded inequality and $\lim c_d^{\text{odd}} = 1/3$ (Part III), the face cover at Proposition 17.7 (Part IV). 64: binomial separation at Gowers’s Section 4 (Part III). 65 and 69: Boolean phase integration in every degree (Part IV). 66: partition normal forms and the fourth energy level (Part II). 67: unequal budgets (Part I), characteristic-three rigidity and the interpolation count (Part II), norm layers (Part III). 68: the energy gap, the coset range and punctures, and exponent 53 (Part II).

Where the eleventh write had to choose. Each delivered section is printed whole in exactly one Part, by its main subject; integration, verification, question and conclusion sections and the appendices are in Part V, as before. Sources 60, 63, 67 and 68 are split by subject (front-matter notes give the map). Where sources prove the same theorem: sources 65 and 69 prove the all-degree Boolean theorem independently; source 65’s statement also covers $d = 3$ and carries the radical-sensitive cap, so it is printed first and source 69 is a credited second route printed whole with its further results, as sources 46 and 43 were; source 63’s limit supersedes source 60’s bracket, so source 63 opens the Part III chapter and source 60’s Section 2 follows as a second route with its connected-limit theorem. In the restriction chapter source 68 comes first (the best exponent for arbitrary targets, continuing source 58), then source 62 (the best powers for cyclic targets), source 61 (exact restriction on arbitrary groups) and source 60’s prime-modulus Section 9; a note compares all of them. In the energy chapter source 68’s universal inequalities come before source 66’s ordered classification.

F.5 Relation to the formal project

The formal project consists of the corrected transcription (`Combinatorics/Ramsey/Papers/sz-thm-gowers-proof`), the Lean development (`Combinatorics/Ramsey/Lean/GowersSzemerédi`, namespace `LeanProofs.GowersSzemerédi`) and the ledger `Combinatorics/Ramsey/gowers-proof-status.json`, which records each numbered statement of the paper as a Lean proposition and its exact companion theorem if there is one. The sources compare their results with *corrected* statements of this project: Lemma 2.3 with denominator 16, Corollary 2.4 with 128π and $\alpha/2$, Corollary 2.5’s exported pair and the stronger pair inside its proof, and so on.

Ledger at the write. At HEAD 25df8755d (6 October 2026) the ledger records 98 statements with exact companions and 22 open, against 88 and 32 at all three pins. The ten statements closed since the pins are `corollary_5_6`, `corollary_5_7`, `lemma_5_9`, `lemma_5_10`, `corollary_5_11`, `lemma_5_14`, `lemma_16_1`, `lemma_16_3`, `lemma_16_5` and `lemma_16_6`. The development is being extended by another session as this report is written, so these counts may already have moved; the crosswalk (Section V.1) records every declaration it cites with file and line at 25df8755d. Two stale statements of the sources are corrected in dated notes: source 02’s description of Corollary 5.6 as an unverified import (it is now `corollary_5_6_holds`), and source 03’s (and implicitly source 01’s) count of 88 exact and 32 open statements.

What is formalized. One statement of the five sources is formalized: source 04’s Lemma III.4.5 (lifting a modular progression from a short interval), with source 03’s rectification remark, as `hasNatAP_of_short_modular_sequence` and `hasNatAP_of_hasModAP_image` (`Proofs18IntervalTransfer.lean`, lines 22 and 92). Those declarations were added in 327773149, source 04’s own pin, before the sources arrived. Some baselines the sources improve are proof-level facts of the development: the stronger Corollary 2.5 pair (`cor25_large_scale`) and the exact second moment of Lemma 9.2 (inside `lemma_9_2_holds`). The Gowers–Cauchy–Schwarz inequality (`gowers_cauchy_schwarz`) and the triangle inequality (`gowersNorm_add_le`) are formal tools that sources 03 and 04 use. The simultaneous partition of sources 01 and 05 is not the formal degree-one

partition `section5_strong_linear_partition` (`Proofs05LinearPartition.lean`, line 86), which is a one-phase statement with target lengths up to $r^{1/16}$ and diameter $r^{-1/8}N$. *Nothing else in this report is formalized, and the report gains no formal status* by its placement beside the Lean development. Each source’s formalization plan is printed in Part V.

[write] 6 October 2026; sources 06–10. *Sources 06–10.* At HEAD `764740f07` the ledger records 99 exact companions and 21 open statements (crosswalk, Section V.1). No statement of sources 06–10 is formalized. Two of them strengthen the conclusion of a formal statement under exactly its hypotheses: source 07’s Corollary II.17.3 (order five in place of order two in `lemma_10_9`) and, by a short argument recorded at the write, source 10’s Corollary II.21.2 and Theorem II.21.1 (which imply the conclusions of `corollary_7_6` and `lemma_7_5`); none of these implications is formalized. Source 07’s order nine needs hypotheses that `lemma_10_9` does not have. The report still gains no formal status by its placement.

[write] 6 October 2026; sources 11–14. *Sources 11–14.* At HEAD `e1db00e99` the ledger records 100 exact companions and 20 open statements, and `FORMALIZATION_STATUS.txt` 206 modules and 1,879 audited theorems (since the second write `lemma_13_8` has closed); at the pins of sources 11–14 (`40f885098`, `9ea383e9b`) it recorded 97 and 23 (154 modules, 1,508 theorems); `lemma_16_6`, `lemma_16_9` and `lemma_13_8` have closed since. No statement of sources 11–14 is formalized, and none gains formal status by its placement. The formal inputs they rely on are named in the crosswalk (Section V.1): source 12’s component input (its Lemma II.40.1) is the output of `lemma_10_3_holds` through `lemma_10_9_holds`, its Appendix A matches `lemma_10_6_holds` with `extractionMode`, and its scalar linearity follows from `corollary_10_14_holds`; source 13’s rectification remark is the already formal `hasNatAP_of_short_modular_sequence`; source 11’s Corollary II.32.2 consumes `proposition_6_1`. One more baseline is a proof-level fact of the development: the proof of `theorem_10_13_holds` already establishes the retained fraction $5\alpha^5/98304$ before exporting $\alpha^6/20000$ (note in Section F.17).

[write] 6 October 2026; sources 15–20. *Sources 15–20.* At HEAD `0b4890a1c` the ledger records 101 exact companions and 19 open statements, and `FORMALIZATION_STATUS.txt` 233 modules and 2,138 audited public theorems; since the third write `lemma_13_7` has closed. At the pins it recorded 97 and 23 (`9ea383e9b`), 98 and 22 (`dca9f0638`) and 99 and 21 (`1bd57e91c`). No statement of sources 15–20 is formalized, and none gains formal status by its placement. Four imply conclusions of formal statements with better constants, none of them formally (crosswalk, Section V.1): source 16’s Theorem F.21.1 implies `proposition_17_7`; source 17’s theorem implies `lemma_9_3` and `corollary_9_4` for prime N ; source 18’s Theorem II.66.8 implies the bound of `lemma_15_4` for $d \geq 2$, $p > 2d$, and its Corollary II.66.9 would sharpen an internal constant of `lemma_15_5_holds`. Three sources understate the formal side: the unweakened Corollary 14.6 exponent is the private, kernel-checked `cor146_direct_lower` (source 15); Proposition 17.7 is the exact companion `proposition_17_7_holds` (source 16); the all- N Lemma 9.3 and Corollary 9.4 are exact companions (source 17). Since source 19’s pin the formal Section 16 has gained results that confirm its audit points and weaker analogues of its local lift (note after its Corollary II.70.8); `lemma_16_10` is still open.

[write] 6 October 2026; sources 21–27. *Sources 21–27.* At HEAD `2dfe46f44` the ledger records 101 exact companions and 19 open statements, and `FORMALIZATION_STATUS.txt` 267 modules and 2,424 audited public theorems; nothing has closed since the fourth write. At the pins it recorded 99 and 21 (`43306653a`, `65f4e039f`, `568d1802e`, `c94998a5d`) and 97 and 23 (`2e0553b0c`). No statement of sources 21–27 is formalized, and none gains formal status by its placement. Two are formalization candidates, named as such and claimed by nobody: source 26’s Corollary I.41.3 gives a written proof of the degree-one instance of the open `corollary_5_8` (`Section05.lean`, line 242), whose ledger status does not change, because the statement quantifies over every degree; and source 23’s common base partition (Theorem II.103.6) would remove the square-root width loss of the formal `section16_local_affine_lift_sqrt`, provided the formal Section 16 interface accepts its long final axes, which is not checked here. One source overstates a gap: source 24’s Proposition 7.3 “in every abelian group” is already the exact companion `proposition_7_3_group_holds` (note in Section F.29). The sharper bounds of sources 22, 25 and 27 at Lemma 15.4 imply the formal bound of `lemma_15_4` in their ranges, as source 18’s does; none is formalized. Source 21’s theorem concerns `lemma_3_10` only through its real centered setting; the formal statement is unchanged.

[write] **6 October 2026; sources 28–33.** *Sources 28–33.* At HEAD 5d4718772 the ledger records 105 exact companions and 15 open statements, and `FORMALIZATION_STATUS.txt` 297 modules and 2,676 audited public theorems; since the fifth write (HEAD 2dfe46f44) `theorem_7_1`, `theorem_7_2`, `theorem_8_2` and `corollary_8_3` have closed, none of them a statement that sources 28–33 compare with. At the pins it recorded 99 and 21 (c94998a5d, e3125e698, 9f83dbe0f) and 100 and 20 (8d65d3c9f). No statement of sources 28–33 is formalized, and none gains formal status by its placement. Formalization candidates, named as such and claimed by nobody: source 32’s Theorem F.37.1, with source 26’s corollary, gives written proofs of the instances $k \leq 2$ of the open `corollary_5_8` (`Section05.lean`, line 242), whose ledger status does not change; source 33’s linear branch strengthens the conclusion of the exact `lemma_12_3` (`Sections12_13.lean`, line 151), subject to matching the two counts’ conventions. Related exact statements (unchanged): `theorem_10_13_holds`, `lemma_10_10_holds`, `corollary_10_11_holds`, `lemma_10_12_holds` (sources 28, 30), `proposition_6_1_holds`, `lemma_7_5_holds`, `corollary_7_6_holds` (sources 29, 31), `lemma_9_3_holds` (source 29, whose theorem is for fixed fields and does not imply it), `lemma_12_3_holds` (source 33) and `lemma_3_10_holds` (sources 28, 29); no declaration concerns c_d^{odd} for $d \geq 4$.

[write] **6 October 2026; sources 34–39.** *Sources 34–39.* At HEAD 89c1cbc38 the ledger records 107 exact companions and 13 open statements, and `FORMALIZATION_STATUS.txt` 410 modules and 2,840 audited public theorems; since the sixth write `lemma_16_4` (c450a0c73) and `theorem_1_2` (5ea1d06dd, from a vendored hypergraph-removal theorem) have closed, neither a statement that sources 34–39 compare with. At the pins of sources 34–39 it recorded 100 and 20. No statement of sources 34–39 is formalized, and none gains formal status by its placement. Two exact formal statements are now also implied by written results: `lemma_5_2` by source 37’s Corollary I.55.2, and `lemma_5_5`, for $t \geq \text{weylThreshold } k$ and with the formal `lemma_5_3` as input, by its Corollary I.57.3 (proved at the write; notes there); both already have exact companions, so nothing changes. Formalization candidates, named as such and claimed by nobody: source 37’s lower-bound certificate for Lemma 5.2; source 39’s sampler at the `Section16SampledOrAnchoredOn` interface (source 19’s first), for the open `lemma_16_10`; source 36’s bilinear estimate for rank-two forms; source 34’s exact reduction, small-prime certificates and fourth-norm inequality. A marked question records the intake’s observation that source 35’s Corollary II.145.3 gives `lemma_15_5` with an explicit threshold (item II.50). Related exact statements (unchanged): `lemma_3_10_holds` (sources 34, 35, 38), `lemma_5_2_holds`, `lemma_5_3_holds`, `lemma_5_5_holds` (source 37), `lemma_5_9_holds` (source 35), `proposition_6_1_holds` (source 38), `lemma_10_9_holds` (source 34), `lemma_14_8_holds`, `lemma_15_5_holds` (source 35), `proposition_17_2_holds`, `proposition_17_7_holds` (source 36); `lemma_16_10` (source 39) is open.

[write] **7 October 2026; sources 40–46.** *Sources 40–46.* At HEAD 6d4942222 the ledger records 113 exact companions and 7 open statements, and `FORMALIZATION_STATUS.txt` 713 modules and 4,583 audited public theorems; since the seventh write `corollary_5_8`, `lemma_13_5`, `lemma_13_6`, `lemma_13_9`, `corollary_13_10` and `theorem_13_12` have closed. *corollary_5_8 is now an exact companion:* `corollary_5_8_holds` (`Proofs05ThresholdFreeLocalization.lean`, line 86), committed in b548d998e on 6 October 2026 (22:15), while this write was in progress. The formal proof follows the two-regime architecture of sources 26, 32 and 41 (Section V.1): the variance lemma (`exists_proper_progression_density_increment`, `Proofs05VarianceIncrement.lean`, line 89; source 41’s Lemma I.66.1, which is also source 26’s and source 32’s), the cover-mass identities, the phase transfer, the high-correlation assembly, the reduction `corollary_5_8_holds_of_threshold_free_localization` (`Proofs05ThresholdFreeCorollary.lean`, line 145) to the localization property `ThresholdFreePolynomialLocalization` (line 16), and that property in every degree (`thresholdFreePolynomialLocalization_all`, `Proofs05ThresholdFreeLocalization.lean`, line 70), proved with the repository’s own finite Weyl recurrence and budgets rather than source 41’s recurrence (`FORMALIZATION_STATUS.txt`: source 41 “gives a related stronger written localization budget; this formal proof uses the already checked finite Weyl estimate”). It proves the catalogue statement $(\text{length } (N/M)^{1/K_k}/8, \text{ gain } \alpha/16)$; source 41’s own Theorem F.46.1 (exponent $1/D_k$, gain $\alpha/3$), its recurrence and its budgets are not formalized. Source 41’s Theorem I.65.3 with Lemma I.65.4, and Lemma I.65.1 at degree one, give the same localization property on paper (checked at the write; note after Lemma I.65.4). No statement of sources 40, 42–46 is formalized either, and none gains formal status by its placement. Related exact statements: `corollary_5_8_holds` (sources 41 and 42, above; new), `lemma_15_4_holds` (source 44), `lemma_15_5_holds` (source 40), `lemma_3_10_holds` (source 45), `theorem_7_2_holds` (source 45’s energy

material; closed before this write), `lemma_17_1_holds` and `proposition_17_2_holds` (sources 43 and 46, whose results are over $\mathbb{F}_2^n$, outside the formal setting). Formalization candidates, named as such and claimed by nobody: source 41's stronger theorem, with the exponent $1/D_k$ and the gain $\alpha/3$, along its recurrence and budgets (the intake recommended its route for `corollary_5_8`, which the repository has since closed by its own); source 42's arithmetic increment in place of the $N \geq L^3$ variance lemma; source 44's width lemma and signed-collision product; source 45's ordered-completion identity and rational U^4 certificate; source 43's orthogonality theorem and Boolean primitive formula.

[write] **7 October 2026; sources 47–53.** *Sources 47–53.* At HEAD 06b30ca52 the ledger records 113 exact companions and 7 open statements (unchanged since the eighth write's HEAD), and `FORMALIZATION_STATUS.txt` 746 modules and 4,742 audited public theorems. At the pins of sources 47–53 it recorded 101 and 19; closed since then: `corollary_5_8`, `theorem_13_12`, `lemma_13_5`, `lemma_13_6`, `lemma_13_9`, `corollary_13_10`, `theorem_1_2`, `theorem_7_1`, `theorem_7_2`, `theorem_8_2`, `corollary_8_3` and `lemma_16_4`. *Formal work on the subjects of these sources since their pins* (Section V.1): the negation of `lemma_16_10`, along source 50's route (`lemma_16_10_packaged_premises_counterexample`, `Proofs16PackagedPremisesCounterexample.lean`, line 13), with a positive explicit cover interface and the endpoint characterization of source 50's Proposition II.234.1; the ledger keeps `lemma_16_10` open and preserves the false predicate pending a faithful repair. A repaired Section 13 chain in which `corollary_13_10` consumes construction geometry (so source 51's obstruction concerns only the preserved older interface) and an explicit density-uniform threshold for square extraction (`section13_complete_square_extraction_explicit`, `Proofs13ExplicitGeometricThreshold.lean`, line 24), with the side $N^{\gamma(\delta)}$ at the lower density; source 53's closed threshold, actual-density side $N^{3\gamma(a)}$ and mass are not formalized. The two Stage 13.7 encodings refuted by source 52 are preserved and unproved, and their negations are not formalized. No statement of sources 47, 48 and 49 is formalized, and none of sources 47–53 gains formal status by its placement. Related exact statements: `lemma_9_3_holds` (source 48), `lemma_15_4_holds` and `lemma_3_10_holds` (source 49), `lemma_17_1_holds` and `proposition_17_2_holds` (source 47, whose results are over $\mathbb{F}_2^n$, outside the formal setting), `lemma_13_7_holds`, `corollary_13_10_holds` and `theorem_13_12_holds` (sources 51–53). Formalization candidates, named as such and claimed by nobody: source 52's two negations (the modular unit lemma, the Freiman no-wrap argument, the arrangement count and one real-power inequality, as it lists); source 53's closed threshold and actual-density exponent, as a strengthening of the uniform theorem, with the exported length $\ell R \leq m + 1$; source 49's translation reduction, rank-two classification and finite certificates; source 47's even-period lemma and canonical induction; source 48's inventory invariants.

[write] **7 October 2026; sources 54–59.** *Sources 54–59.* At HEAD fe7165a3c the ledger records 113 exact companions and 7 open statements (unchanged since the ninth write), and `FORMALIZATION_STATUS.txt` 755 modules and 4,836 audited public theorems. At the pins of sources 54–59 it recorded 103 and 17 (source 54) and 105 and 15 (the others); `theorem_13_12` was open at all of them. *Formal work on the subjects of these sources since their pins:* `theorem_13_12` is an exact companion since d48e8af00 (6 October 2026, 23:55), `theorem_13_12_holds` (`Proofs13EndpointPrintedTheorem.lean`, line 73), *along source 57's route:* the repository's own uniform Fejér purification (e40138bd5), propagated through the prepared frequency graph to density $(\alpha/2)^{4,207,554,485}$ and the length exponent 2^{53} (`theorem_13_12_with_fejer_explicit_bound`, `Proofs13FejerExplicitExponent.lean`, line 55), compared with the printed length for $\alpha \leq 1023/1024$ (`theorem_13_12_away_from_one`, `Proofs13FejerPrintedRange.lean`, line 39), and source 57's near-maximal endpoint for $\alpha > 1023/1024$ (`endpoint_global_fourier`, `Proofs13EndpointGlobalFourier.lean`, line 12; `theorem_13_12_near_one`, `Proofs13EndpointPrintedTheorem.lean`, line 12), crediting source 57 in `FORMALIZATION_STATUS.txt`. The formal density exponent is 2^{42} , not source 57's 2^{41} ; the repository also proves 2^{61} in place of the printed length exponent 2^{70} (`theorem_13_12_with_improved_printed_bounds`, `Proofs13ImprovedPrintedBounds.lean`, line 11), its own strengthening. Notes in Section II.287. No other statement of sources 54–59 is formalized, and none gains formal status by its placement. Related exact statements: `proposition_6_1_holds`, `lemma_6_2_holds` and `corollary_7_6_holds` (sources 55, 56, 57, 59), `lemma_7_5_holds` (59), `proposition_8_1_holds` (54), `lemma_9_3_holds` and `corollary_9_4_holds` (58), `lemma_10_9_holds` (56), `lemma_13_11_holds` and `theorem_13_12_holds` (57). Formalization candidates, named as such and claimed by nobody: source 58's Theorem A as a strengthened `lemma_9_3` (finite combinatorics and five integer certificates); source 54's Theorem F.59.1 beside `proposition_8_1`;

source 55’s explicit local theorem; source 59’s separation ranks and signed affine identities; source 57’s partition classification; source 56’s five-module plan.

[write] **7 October 2026; sources 60–69.** *Sources 60–69.* At HEAD fb9602e55 the ledger records 114 exact companions and 6 open statements (113 and 7 at the tenth write; the change is the repair of Lemma 16.10, whose companion `lemma_16_10_holds` is counted but not yet kernel-checked, Section V.1), and `FORMALIZATION_STATUS.txt` 764 modules and 4,877 audited public theorems in its last complete audit, which predates the repair. At the pins of sources 60–69 (6 October 2026, 21:05–22:07) it recorded 107 and 13. *Formal work on the subjects of these sources since their pins:* the affine case of source 67’s interpolation count is formalized in the repository’s own form (`Section16SampledOrAnchoredOn.compressed_multilinear_cover`, `Proofs16CompressedInterpolation.lean`, line 56, 081abc4d3, 7 October 2026, 08:49: a family of $\max(p, \binom{r}{2}p^2)$ candidates for arbitrary samples with replacement), propagated through the repository’s Section 16 affine lift with the strict improvement over $rb + r^2b^2$ (`section16_compressed_budget_lt_old`, `Proofs16CompressedCandidateBounds.lean`, line 42; `section16UniformLiftGraphBudget_lt_previous`, `Proofs16CompressedLiftImprovement.lean`, line 39); `FORMALIZATION_STATUS.txt` names source 67’s theorem as its origin and records that the higher-degree count and the lower constructions are not formalized; the repaired `lemma_16_10` (`Section16.lean`, line 912) bounds its graph count in the same compressed form, and does not give the printed unit-parameter conclusion. No other statement of sources 60–69 is formalized, and none gains formal status by its placement; nothing in them closes an open ledger statement. Related exact statements: `lemma_5_2_holds` (67), `proposition_6_1_holds` (67), `lemma_7_5_holds` and `corollary_7_6_holds` (61), `lemma_9_3_holds` and `corollary_9_4_holds` (60, 61, 62, 68), `lemma_12_3_holds` (60), `lemma_3_10_holds` (63), `proposition_17_2_holds` (65, 69) and `proposition_17_7_holds` (63); `conjecture_4_1` and `conjecture_4_2` (64) are proposition-valued definitions, not asserted. Formalization candidates, named as such and claimed by nobody: source 68’s exponent-53 theorem and source 61’s threshold-free exact restriction as strengthened `lemma_9_3`; source 68’s energy gap and coset theorem; source 66’s partition normal form; source 65’s integration criterion and endpoint recognition; source 63’s bounded inequality and face-cover identity; source 64’s seed certificate.

F.6 Source 01. The two interfaces and the main conclusions

Source 01, delivered Section 1.

Delivered title block. “Sharp Local Refinements of Gowers’s Szemerédi Argument / Density transfer, carry-pattern entropy, and formalization interfaces / Research report prepared for Vladimir Reshetnikov / Mathematical development and computational checks with ChatGPT / 6 October 2026”; PDF author field “Research report prepared with ChatGPT”. Archive `Gowers_Quantitative_Refinements(1).zip`, manuscript `gowers_refinements.tex` (1,980 lines, 31 pages), pin 3a25020a6. Printed below: the abstract, the status paragraph and Section 1; Sections 2–8 are in Part I, 9–13 in Part IV, 14–16 and Appendices A–B in Part V.

Abstract of source 01 (as delivered). We develop two quantitative refinements of components of Gowers’s 2001 proof. First, a convex-envelope calculation solves the finite weighted extremal problem of converting correlation with a cellwise short-arc phase into a density tail. The resulting bound is sharp in the stated information model, keeps track of occupied mass in exceptional cells, and yields an explicit Fourier density increment on a progression of length on the scale $\sqrt{\eta N / [\delta(1 - \delta)]}$, in place of the $\sqrt{\eta^3 N}$ length in the printed Corollary 2.5. Second, the number of Boolean subset-sum floor patterns with coefficients in $(-r, r)$ is exactly $(2r)^k C_k$, where C_k counts fractional carry patterns. We prove a quadratic-exponent upper bound, an elementary $2^{\lfloor k^2/4 \rfloor}$ lower bound, and, using the classical enumeration of linear threshold functions, $\log_2 C_k \sim k^2$. A polynomial-feature extension has matching leading entropy for fixed degree. An affine version bounds translation classes independently of the modulus, providing a smaller combinatorial loss at the phase-removal interface in Section 17. Exact rational certificates give $C_1 = 1$, $C_2 = 2$, $C_3 = 10$, and $C_4 = 154$. All local results are proved here; the threshold-function input is explicitly

external. These refinements are not asserted to improve the best global bounds for Szemerédi's theorem, and no Lean kernel verification is claimed.

Status and scope. This is a research manuscript, not a refereed publication. “Sharp” below always refers to a precisely specified extremal problem, not to every arithmetic realization of that problem. The work establishes strengthened statements relative to the cited components of Gowers's paper; an exhaustive priority search for every formulation has not been completed. The exact factorization, extremal transfer formulas, and their proofs can be assessed independently of any novelty claim.

Two recurrent losses in a density-increment proof are logically distinct. One converts phase correlation into a local increase of density. The other counts the discrete floor patterns created when a continuous phase is cut into progression cells. Improving one does not automatically improve every subsequent loss. We keep these interfaces separate.

F.6.1 An explicit density-increment replacement

Write $e(t) = \exp(2\pi it)$ and $[N]_0 = \{0, \dots, N-1\}$. Suppose $A \subseteq [N]_0$ has density $0 < \delta < 1$ and, for some real ξ ,

$$\left| \frac{1}{N} \sum_{n=0}^{N-1} (\mathbf{1}_A(n) - \delta) e(\xi n) \right| \geq \eta > 0. \quad (\text{F.6.1})$$

Put $\kappa = \eta/[2\delta(1-\delta)]$. Necessarily $0 < \kappa \leq 1$. Theorem I.6.1 proves that a proper arithmetic progression $P \subseteq [N]_0$ exists with

$$|P| \geq \max \left\{ 1, \left\lfloor \sqrt{\frac{N}{2\pi}} \arcsin \frac{\eta}{4\delta(1-\delta)} \right\rfloor \right\}, \quad (\text{F.6.2})$$

$$\frac{|A \cap P|}{|P|} \geq \delta + \frac{3\delta(1-\delta)\kappa}{4 - (4-3\delta)\kappa} \geq \delta + \frac{3\eta}{8}. \quad (\text{F.6.3})$$

No primality hypothesis is needed. A nonzero cyclic Fourier coefficient is a special case of (F.6.1). The printed Corollary 2.5 of Gowers gives length $\sqrt{\eta^3 N/(128\pi)}$ and increment $\eta/8$ in that special case [1, Corollary 2.5].

The ingredients here are elementary. In particular, the fact that a minimum-cell-size partition avoids some subsequent length loss is already implicit in the preceding Corollary 2.4. Our point is the density-sensitive, integer-exact formulation and the optimal finite weighted transfer theorem behind it, not a claim that the Fourier density-increment strategy itself is new.

F.6.2 A quadratic rather than cubic floor-pattern exponent

For $\alpha \in \mathbb{R}^k$, define

$$F_\alpha(S) = \left\lfloor \sum_{i \in S} \alpha_i \right\rfloor, \quad S \subseteq [k].$$

For integers $r_i \geq 1$, let $\mathcal{F}(\mathbf{r})$ be the set of these functions when $-r_i < \alpha_i < r_i$. Define

$$C_k = \# \left\{ S \mapsto \left\lfloor \sum_{i \in S} x_i \right\rfloor : x \in [0, 1)^k \right\}.$$

Theorem IV.1.2 gives the exact identity

$$|\mathcal{F}(\mathbf{r})| = \left(\prod_{i=1}^k 2r_i \right) C_k. \quad (\text{F.6.4})$$

Let

$$H_k = (k-2)2^{k-1} + 1, \quad \mathcal{R}(h, d) = \sum_{j=0}^d \binom{h}{j},$$

where $\binom{h}{j} = 0$ for $j > h$. We prove

$$2^{\lfloor k^2/4 \rfloor} \leq C_k \leq \mathcal{R}(H_k, k). \quad (\text{F.6.5})$$

In the isotropic case this replaces the 2^{2rk^3} upper bound in Gowers’s Corollary 17.5 by $(2r)^k \mathcal{R}(H_k, k)$ [1, Corollaries 17.5–17.6]. The threshold-function theorem discussed in Section IV.2 then yields

$$\log_2 C_k = k^2 + o(k^2). \quad (\text{F.6.6})$$

Thus a subquadratic exponent cannot bound the full family, even when the coefficient range is fixed. This optimal leading exponent is obtained by combining a reduction proved here with a classical external theorem.

F.6.3 What is, and is not, concluded

The proofs below solve two specific auxiliary problems: the optimal correlation-to-density-tail problem with freely placed short-arc phases, and the exact dependence of floor-pattern counts on integral coefficient ranges. They also locate a smaller translation-class budget for a phase-removal step. They do not supply a new inverse theorem for higher Gowers norms or a new globally optimized iteration.

For comparison, Leng, Sah, and Sawhney prove, for $k \geq 5$ and some $c_k > 0$,

$$r_k(N) \ll N \exp(-(\log \log N)^{c_k}),$$

where $r_k(N)$ is the largest size of a k -term-progression-free subset of $[N]$ [2]. That checked benchmark is already much stronger than what is obtained by merely adjusting constants in the 2001 argument.

F.7 Source 02. Scope, sources, and the precise improvements

Source 02, delivered Section 1.

Delivered title block. “Energy-Sensitive Inverse Steps and Sharp Density Extraction / Quantitative refinements of Gowers’s Szemerédi argument / Research article prepared for Vladimir Reshetnikov / October 6, 2026”. Archive `Gowers_Szemerédi_Refinements.zip`, manuscript `gowers_refinements.tex` (823 lines, 20 pages; not the same file as source 01’s `gowers_refinements.tex`). No pin in its README; its Appendix A records the observed branch reference `66c02f8a1`. Printed below: the abstract and Section 1; Sections 2–5 are in Part II, 6–9 and Appendix B in Part I, 10–12 and Appendix A in Part V.

Abstract of source 02 (as delivered). We strengthen two interfaces in the proof of Szemerédi’s theorem: the passage from derivative Fourier correlations to additive energy, and

the passage from local phase correlations to a density increment. The first result is a sharp, mixed-function, weighted extension of Gowers’s Proposition 6.1, valid on every finite abelian group without a pointwise boundedness assumption. If $m(f) = \mathbb{E}|f|^2$, if $\mathcal{U}(f) = \|f\|_{U^3}^8$, and if $\mathcal{T}_w(f, g)$ is a weighted derivative-correlation statistic, then

$$\mathcal{T}_w(f, g)^4 \leq \min\{m(f)^4 \mathcal{U}(g), m(g)^4 \mathcal{U}(f)\} \mathcal{E}(w).$$

Here $\mathcal{E}(w)$ is explicitly normalized additive energy on $G \times \widehat{G}$. We prove the identities underlying this inequality, give equality examples, and quantify its two Cauchy–Schwarz deficits. For balanced indicators, the resulting graph-energy estimate retains the variance $\delta(1 - \delta)$ rather than discarding it.

The second result is a sharp real-valued phase-freezing inequality. If the phases in every cell have chordal diameter at most 2ρ , local correlation C , signed cell mass S , and total absolute mass L satisfy $C \leq \rho L + (1 - \rho)S$. A sharp one-sided moment inequality then gives the total mass of cells with any prescribed positive density increment. This strengthens Lemma 5.15, yields an explicit refinement theorem using a single global phase budget, and permits closed-form optimization of that budget. The elementary results are proved in full; the polynomial partition specialization explicitly imports the repository’s corrected Corollary 5.6. These are local quantitative refinements, not a claimed improvement of the best global bound for Szemerédi’s theorem. Literature priority for the exact combined formulations is not established.

F.7.1 The two proof interfaces

Gowers’s proof [1] repeatedly turns a failure of uniformity into structure and then into increased density on an arithmetic progression. The supplied ProveIt transcription and Lean development [20] make the intermediate interfaces particularly useful targets: they have concrete finite statements, explicit normalizations, and independently reusable conclusions.

The starting point for our first improvement is Proposition 6.1. In the paper’s unnormalized Fourier convention, if $f : \mathbb{Z}_N \rightarrow \mathbb{C}$ satisfies $|f| \leq 1$, $B \subseteq \mathbb{Z}_N$, and $\phi : B \rightarrow \mathbb{Z}_N$ satisfies

$$\sum_{h \in B} |\widehat{\Delta_h f}(\phi(h))|^2 \geq \alpha N^3, \tag{F.7.1}$$

then the graph of ϕ contains at least $\alpha^4 N^3$ additive quadruples. The proof uses Cauchy–Schwarz and Fourier orthogonality, but bounds two quantities depending on f by their worst-case values. We retain both quantities and allow a weighted, multivalued spectrum and two different input functions.

For our second improvement, Lemma 5.15 assumes a real, mean-zero, 1-bounded function and a partition $P_1, \dots, P_M$ such that

$$\sum_j \left| \sum_{x \in P_j} f(x) \right| \geq \alpha N.$$

It produces a cell with mean at least $\alpha/4$ and size at least $\alpha N/(4M)$. We obtain a full tradeoff between the requested cell mean and the mass of all cells attaining it. The estimate uses the asymmetric range of a balanced indicator, $[-\delta, 1 - \delta]$, instead of only $[-1, 1]$.

Between these two interfaces lie the phase partitions of Section 5. We sharpen the geometric step that removes a nearly constant phase and then track one global error budget through an arbitrary collection of input progressions. No lower bound on each individual local correlation is needed.

F.7.2 What is and is not claimed

The weighted Fourier-energy inequality, the real-phase inequalities, the sharp moment bound, and their stated consequences have complete proofs below. The coefficient in the main energy inequality and the constants in the abstract phase and moment inequalities are optimal in the senses explicitly proved. “Sharp” does not mean that every later progression-length bound, or the complete density-increment iteration, is optimal.

The estimate $\|f\|_{U^3} \leq \|f\|_2$ and the general subject of extremizers are established mathematics; the broader near-extremizer theory is treated by Eisner and Tao [7]. Our elementary proof of that estimate is included to make the constants and the treatment of torsion transparent. Similarly, we do not claim that retaining a Cauchy–Schwarz denominator is a new general technique.

These results do not supply a new best asymptotic estimate for $r_k(N)$. For comparison, Leng, Sah, and Sawhney prove, for $k \geq 5$, a bound of the form

$$r_k(N) \ll N \exp(-(\log \log N)^{c_k}), \quad c_k > 0,$$

using quasipolynomial inverse-theorem bounds and a different quantitative organization of the density-increment argument [2, 3]. We use that published preprint statement as a benchmark, not as a claim to have exhaustively determined the latest record in every progression length.

The exact combined formulations here are independently derived refinements of the supplied proof. The literature check does not establish publication priority, and no previously documented major open problem is declared solved. Section V.9 separates useful next research problems from results proved here.

F.8 Source 03. Scope, starting point, and principal results

Source 03, delivered Section 1.

Delivered title block. “Additive Combinatorics / Sharper local estimates in Gowers’s proof of Szemerédi’s theorem / Density increments, exact cube expansions, and dependent random choice / Research report prepared for Vladimir Reshetnikov / 6 October 2026”. Archive `gowers_local_refinements.zip`, manuscript `gowers_local_refinements.tex` (2,392 lines, 36 pages), pin 3a25020a6. Its title page also says: “Source basis: Gowers (2001) and the two requested `ProveIt` directories, pinned to the revision documented in Section V.13. The accompanying code checks finite instances; the proofs in the text establish the general statements.” Printed below: the abstract and Section 1; Sections 2–4 are in Part I, 5–6 in Part III, 7–8 in Part II, 9–11 and Appendix A in Part V.

Abstract of source 03 (as delivered). We develop explicit local strengthenings of several estimates used in Gowers’s proof of Szemerédi’s theorem and relate them to the accompanying `ProveIt` Lean development. First, an exact asymmetric selection inequality converts partition discrepancy into a sharp lower bound for the mass of cells above a chosen density threshold. Combined with an elementary progression partition, it gives a Fourier density increment whose length retains a factor $[\delta(1 - \delta)]^{-1/2}$. Second, a telescoping progression count replaces an exponential error coefficient by a density-weighted sum. Third, we compute the entire fourth-order term of the centered cube count on every finite abelian group. Its two contributions are ordinary Fourier energy and an explicit correction from the quotient by doubling. The resulting fourth-order error and its universal leading constant are sharp, including for sequences of actual sets. Finally, a cubic dependent-random-choice certificate improves the exact-threshold subset size in Lemma 7.4 from order $\delta^5 n$ to order $\delta^3 n$; a separate one-witness argument greatly improves the elementary constants in Proposition 7.3. All claimed refinements have complete proofs. We

identify their relation to established methods, distinguish them from stronger modern theorems, and give a concrete formalization and research agenda. No improvement to the global record for Szemerédi’s theorem or solution of a major outstanding conjecture is asserted.

The aim of this report is to strengthen concrete parts of the proof in [1], with statements precise enough to be independently checked and incorporated into the supplied formal development [21]. The two supplied directories have different roles. The paper directory contains an edited transcription of the original article, including finite-scale repairs and explanatory comments. The Lean directory contains a catalogue of numbered propositions together with many companion proofs. We compare both with the published article rather than silently identifying their statements.

The work concentrates on three places where numerical losses can be removed without inventing a new inverse theorem: selection from a partition, expansion of a multilinear count, and extraction of a structured subfamily by random sampling. These are useful mathematical components in their own right. Their improvements are also easier to audit than a claim to have sharpened the final tower of parameters after an unexamined substitution into a long proof.

F.8.1 What is established

The following are the central conclusions. The conventions and all finite-scale assumptions appear in the cited statements within this report.

1. **Optimal selection from asymmetric discrepancy data.** For a mean-zero random variable $X \in [-a, b]$, its positive first moment T determines the exact best universal lower bound

$$\Pr(X > t) \geq \max \left\{ 0, \frac{(a+t)T - at}{a(b-t)} \right\}.$$

Theorem I.14.1 includes extremizers. Specializing to cell averages strengthens Lemma 5.15 and allows the density and size requirements to be chosen simultaneously.

2. **A Fourier increment with the density factor preserved.** For a normalized nonzero Fourier coefficient of size ρ , an explicit proper integer progression has the bounds in (I.15.3). The construction works for composite moduli and handles singletons and integer rounding directly. It improves on the stronger estimate already present inside the repository’s proof, not merely its weaker exported conclusion.
3. **A smaller progression-count error.** Theorem III.1.2 gives a density-weighted sum with $k - 2$ terms. Under the hypotheses of Corollary 3.3, in its original index order, this replaces $2^k \alpha$ by at most $(k - 2)\alpha$. Theorem III.1.2 also permits any family of slopes whose pairwise differences act invertibly on the group.
4. **An exact fourth-order cube expansion.** Theorem III.2.6 computes the leading possible error rather than only bounding it. There are no terms of degrees one, two, or three. Among degree-four supports, only ordinary parallelograms and a second family visible to order-two characters contribute. Corollary III.2.7 changes the displayed error scale in Lemma 3.10 from $\alpha^{1/2^d}$ to $\alpha^{4/2^d}$. The fourth power and the universal leading coefficient are optimal in the precise senses proved below.
5. **A larger common-neighborhood core at the same threshold.** Theorem II.5.2 replaces $2^{-1/2} \delta^5 n$ by $7\delta^3 n / (3\sqrt{6})$, while retaining both the threshold $\delta^2 m / 2$ and the 90% requirement in Lemma 7.4. The certificate also gives weighted higher-tuple versions.

6. **An explicit elementary energy-to-structure improvement.** For $E(A) \geq c|A|^3$, Theorem II.6.3 gives

$$|A'| \geq (c/8)|A|, \quad |A' - A'| \leq 2^{19}c^{-9}|A|.$$

Both exponents improve on Proposition 7.3’s displayed constants. The proof works in every abelian group and uses only elementary counting. It is substantially weaker than modern BSG results.

Several ingredients here are established general methods: normalized Gowers norms, telescoping, dependent random choice, and the passage from many additive quadruples to a small difference set. The contribution of this report is the explicitly proved package of refinements, the exact quartic classification with its sharpness statements, the optimized scalar selection problems, and a source-aware route to formalization. We make no historical priority claim for these formulations without a separate, exhaustive literature investigation.

F.8.2 Relation to contemporary quantitative results

Let $r_k(N)$ denote the largest size of a subset of $\{1, \dots, N\}$ containing no nonconstant k -term arithmetic progression. Gowers’s quantitative theorem gave a bound of the shape $r_k(N) \ll_k N(\log \log N)^{-c_k}$. Leng, Sah, and Sawhney prove the stronger estimate

$$r_k(N) \ll_k N \exp(-(\log \log N)^{c_k}) \quad (k \geq 5) \quad (\text{F.8.1})$$

for suitable $c_k > 0$ [2]. Their argument uses quasipolynomial inverse-theorem bounds [3] and a density increment strategy related to [5]. None of the estimates established here improves (F.8.1).

The historical questions at the end of Gowers’s paper must also be read in their date. For example, Jaber, Liu, Lovett, Ostuni, and Sawhney prove that a corner-free set in $G \times G$, with G finite abelian, has size at most

$$|G|^2 \exp(-(\log |G|)^{\Omega(1)})$$

[13]. Thus simply restating the 2001 request for a reasonable two-dimensional corners bound would not identify an unsolved problem. Our research questions in Section V.14 instead specify which finer optimization, extension, or formalization remains to be done.

F.9 Source 04. Purpose, results, and scope

Source 04, delivered Section 1.

Delivered title block. “Additive Combinatorics / Sharper Local Bounds in Gowers’s Proof of Szemerédi’s Theorem / Polynomial homomorphism restriction, sharp cube estimates, and floor-pattern complexity / Prepared for Vladimir Reshetnikov / 6 October 2026”. Archive `gowers_quantitative_refinements.zip`, manuscript `gowers_quantitative_refinements.tex` (2,987 lines, 40 pages; the base of the merge, staged as `article.tex`), pin 327773149. Its first lines read “Complete manuscript source. See README.txt for build instructions. Prepared on 2026-10-06. No new Lean kernel verification is claimed.”, and its title page “Companion materials: editable L^AT_EX source, compiled PDF, exact rational pattern certificates, verification programs, and a source manifest.” Printed below: the abstract and Section 1; Sections 2–5 are in Part II, 6–7 in Part III, 8 in Part IV, 9–11 and Appendices A–B in Part V.

Abstract of source 04 (as delivered). We give explicit refinements of local estimates in Gowers’s proof of Szemerédi’s theorem and identify interfaces for their formalization. Exact Fourier second moments sharpen the lifting of respected additive relations. A powered-cosine selection argument, with Fourier-weighted control of exceptional relations and separate treatment of repetitions, gives a polynomial restriction theorem at every fixed tuple order. For real centered functions on odd-order groups, we prove the sharp gap $\|f\|_{U^3}^8 \geq 2\|f\|_{U^2}^8$, classify equality, and obtain stability. The resulting three-cube bound has sharp leading coefficient $6\sqrt{2}$. Without 2- or 3-torsion, its first correction has optimal exponent $16/3$, witnessed by an explicit two-frequency family. Further results retain density factors in progression counting and determine the quadratic exponential scale of fractional floor patterns, improving the associated local phase-extraction loss. The article includes complete proofs, exact examples, supplementary verification programs, a map to the inspected Lean files, and twelve research questions. Improvements are relative to the specified local source estimates; worldwide priority, a new best global Szemerédi bound, and new kernel-checked Lean proofs are not claimed.

This article develops explicit improvements to several local estimates in Gowers’s proof of Szemerédi’s theorem [1]. The starting points are the paper and its companion Lean development in the ProveIt repository [23]. The improvements come from retaining quantities that a convenient coarse estimate discards: the size of a Fourier support measure, cancellation in a centered cube expansion, the simultaneous dependence of floor values on their parameters, and the number of patterns that occur at a fixed parameter.

The most substantial quantitative change concerns the passage from a map respecting many additive quadruples to a restriction respecting almost all additive tuples of a higher order. Here is a concrete consequence, stated at the same order as Gowers’s Corollary 9.4. Suppose that $B \subset \mathbb{F}_p$, $|B| = \beta p$, and $\phi : B \rightarrow \mathbb{F}_p$ respects at least $\gamma|B|^3$ additive quadruples. If $0 < \beta, \gamma, \eta \leq 1$ and

$$p \geq 2^{102} \gamma^{-112} \eta^{-16} \beta^{-1},$$

then there is $B' \subset B$ containing at least

$$2^{-95} \gamma^{112} \eta^{15} |B|^{15}$$

ordered additive 16-tuples, at least a proportion $1 - \eta$ of which are respected by ϕ . Both the exponents and the size hypothesis are explicit. The proof uses one powered cosine filter and bounds the exceptional relations with their Fourier weights. It generalizes to every fixed tuple order $2d$, with a retention bound proportional to $\alpha^{2d} \eta^{2d-1}$ when the initial respected-tuple proportion is normalized by $|B|^{2d-1}$.

The second set of results concerns additive cubes. For a real centered function on a finite abelian group of odd order, we prove the sharp inequality

$$\|f\|_{U^3}^8 \geq 2\|f\|_{U^2}^8.$$

Equality, apart from the zero function, is characterized by support on one conjugate pair of Fourier characters. A quantitative stability statement identifies an almost extremizing Fourier pair. Combined with an exact enumeration of the surviving terms in a centered three-cube expansion, this gives a sharp leading coefficient $6\sqrt{2}$ for the cube excess as a function of $\|f\|_{U^3}$. In groups whose order is coprime to 6, a further refinement replaces a fifth-order remainder by a remainder of order $u^{16/3}$. An explicit two-frequency family shows why that fractional exponent occurs.

The remaining results provide useful quantitative interfaces for the larger proof. A telescoping expansion reduces the exponential counting factor in Corollary 3.3 to a density-weighted sum with

$k - 2$ terms. The statement extends to composite cyclic groups and other finite abelian groups under exact invertibility hypotheses. In Section 17, separating integer parts from fractional parts gives polynomial dependence on the slope range. Matching quadratic-order upper and lower bounds determine the exponential scale of the fractional floor-pattern count. A distinction between global patterns and the patterns for a fixed slope improves a subsequent Cauchy–Schwarz loss.

F.9.1 What is established, and what is not claimed

All results asserted as theorems below have proofs. The supplementary programs check finite instances, exact support counts, coefficient inequalities, and rational certificates. These computations are additional evidence and reproducibility aids; the general mathematical statements do not rest on numerical experimentation.

The claim of improvement is relative to the explicitly identified estimates in Gowers’s paper and the inspected repository files. Hölder interpolation, generalized von Neumann inequalities, and hyperplane-arrangement estimates are established tools. A targeted literature review did not determine priority for every exact refinement assembled here. Accordingly this article does not present the collection as a certified first discovery in the literature.

There is also a material distinction between improving a local lemma and improving the best bound in Szemerédi’s theorem. Leng, Sah, and Sawhney prove, for fixed $k \geq 5$, a bound of the form

$$r_k(N) \ll_k N \exp(-(\log \log N)^{c_k}), \quad c_k > 0,$$

using their quantitative inverse theorem [2, 3]. The local arguments here do not improve that result. Green and Tao’s four-term work [6] provides an additional example of why approximate-homomorphism restriction is a useful component of a density increment proof. Manners’s polynomial true-complexity theorem [14] places the counting arguments in a broader established theory. No unresolved global Szemerédi or inverse-theorem conjecture is declared solved in this article.

F.9.2 The inspected repository

The repository was inspected on 6 October 2026. Its paper transcription already contains editorial corrections, and the Lean directory separates formal statement definitions from companion proof modules. The referenced files agree with the snapshot

327773149bd272239f1f6c0015f76d5097512c88.

In particular, the relevant files include `Proofs09Moments.lean`, `Proofs03CubeUpper.lean`, `Proofs03ProgressionAverage.lean`, and `Proofs03ProgressionCount.lean`. The appendix gives a concrete plan for adding the strengthened results. No new Lean proof is claimed to have been compiled or kernel checked in the present project, and no repository changes are part of this deliverable.

F.10 Source 05. Scope, baseline, and main contributions

Source 05, delivered Section 1.

Delivered title block. “Sharp Finite Phase Partitions and Density-Sensitive Increments / Strengthenings of Gowers’ localization lemmas, simultaneous phases, and spectral-energy extraction / Research note prepared for Vladimir Reshetnikov / 6 October 2026”; empty PDF author field. Archive `gowers_strengthening.zip`, manuscript `article.tex` (847 lines, 18 pages), pin 327773149. Printed below: the abstract and Section 1; Sections 2–6 are in Part I, 7–10 in Part V.

Abstract of source 05 (as delivered). We investigate the finite localization step in Gowers’ proof of Szemerédi’s theorem, using the corrected paper and the corresponding Lean modules in the ProveIt repository as the comparison baseline. Three improvements are established with complete proofs. First, the corrected lower constant $1/4$ in Lemma 2.3 can be replaced by $1/3$ without changing its upper length bound; $1/3$ is the largest possible universal constant in that formulation. Second, a nonzero Fourier coefficient of normalized magnitude at least α for a set of density δ produces an arithmetic progression of length at least

$$\sqrt{\frac{\alpha N}{16\pi\delta(1-\delta)}} \quad \text{and density at least} \quad \delta + \frac{\delta\alpha}{4\delta - \alpha}.$$

An exact integer formula is stronger, and a positive proportion of the partition has a controlled increment. Third, simultaneous localization of m real linear phases has optimal order $1 + (\varepsilon^m N)^{1/(m+1)}$, including a partition with a lower bound on every cell and an explicit algebraic obstruction to a better exponent. A density-sensitive spectral-energy inequality links these partitions to multiplicative increments, and unequal phase tolerances can be optimized explicitly. The methods are elementary refinements of classical arguments; claims of literature priority, a new global Szemerédi bound, and completed Lean verification are not made. Exact-arithmetic partition tests and high-precision Fourier diagnostics accompany the article.

F.10.1 What is being strengthened

Gowers’ proof [1] separates a counting argument from the extraction of structure and then a density increment. Our subject is the last conversion: turning a structured correlation into density on a sufficiently long ordinary arithmetic progression. Even at the linear-phase level, this conversion involves three distinct sources of loss: integer rounding in a partition, the error from replacing an oscillating phase by a constant, and the passage from absolute block sums to a positive block average.

The requested repository contains an edited transcription of the paper and a substantial Lean development [24, 25, 26]. The transcription is not identical to the printed paper. In particular, it explicitly repairs finite rounding issues in Lemma 2.3 and Corollary 2.4. The relevant comparison is therefore with the *corrected repository statements*, not merely with the historical printed constants.

Write $t = rs/M$, where r is the length of the domain interval, M the modulus, and s the allowed image diameter. The repository’s version of Lemma 2.3 asserts a partition into progressions with lengths between $\sqrt{t}/4$ and $\sqrt{t}$. The printed lower constant $1/2$ cannot hold uniformly: when $r = s = M = 5$ and the phase is constant, it would require a partition of five points into sets all of size two. We prove that $1/3$, not $1/4$, is the sharp universal lower constant when the upper constant is fixed at one.

The repository’s Corollary 2.5 states a progression length of at least $\sqrt{\alpha^3 N / (128\pi)}$ and a density increment $\alpha/8$. Its large-scale proof already obtains the stronger pair $\sqrt{\alpha N / (128\pi)}$, $\alpha/4$ before weakening the conclusion [26]. Thus simply noticing this stronger pair would not constitute an advance over the repository. Our density-sensitive length estimate improves beyond that proof-level baseline as well.

F.10.2 A map of the results

The exact finite strengthening is Theorem I.16.4, including the optimality of $1/3$. Corollary I.16.6 simultaneously improves the corrected Corollary 2.4 while keeping its upper length bound. The density-sensitive replacement for Corollary 2.5 is Theorem I.17.4. The simultaneous-phase theorem and its matching exponent obstruction are Theorems I.18.2 and I.18.5. The spectral statement, Theorem I.19.2, depends on the residual-variance estimate

$$E \leq \sqrt{EV} + 2R(v - V), \quad (\text{F.10.1})$$

where E is captured Fourier energy, V is the variance of the block densities, $v = \delta(1 - \delta)$, and R is a weighted phase-oscillation budget.

F.10.3 What is not asserted

The elementary ideas used below—pigeonholing on a torus, grouping residue classes, algebraic norms, and conditional expectation—are established mathematical tools. Our contribution is a set of explicit, proved finite statements and their quantitative combination. We have not established that every displayed refinement is absent from the literature.

Nor do we claim an improvement to the best global bounds for $r_k(N)$. Already in 2024, Leng, Sah, and Sawhney proved, for $k \geq 5$, an estimate of the form $r_k(N) \ll N \exp[-(\log \log N)^{c_k}]$, using substantially deeper inverse and localization results [2]. This is a comparison point, not an assertion of an exhaustive October 2026 literature ranking. Their work also emphasizes simultaneous localization and multiplicative increments. The results here concern an explicit linear-phase interface and do not replace the higher-order part of either proof.

F.11 Source 06. Scope, conclusions, and mathematical provenance

Source 06, delivered Section 1.

Added in the second write, 6 October 2026. Delivered title block. “Sharp Local Bounds in Gowers’ Proof of Szemerédi’s Theorem / Carry patterns, quadratic counting errors, and optimal density extraction / Research report prepared for Vladimir Reshetnikov / 6 October 2026”; the same in the PDF author field. Archive `Gowers_Refinements (1).zip` (an independent manuscript: not an edition of source 10’s `gowers_refinements.zip` or of source 01’s archive), manuscript `Gowers_Refinements.tex` (1,958 lines, 27 pages), pin 7d99e5dee (its `provenance.json`). Printed below: the abstract and Section 1; Section 5 is in Part I, 3 in Part III, 2, 4 and 7 in Part IV, 6, 8 and 9 in Part V.

Abstract of source 06 (as delivered). We develop explicit refinements of several local arguments in Gowers’ proof of Szemerédi’s theorem, with the accompanying ProveIt development as a precise reference point. For Boolean floor patterns with slopes in $(-r, r)^k$, we prove the exact count $(2r)^k C_k$, where C_k is a unit-cube carry count. An elementary arrangement bound and a reduction from Boolean threshold functions show $\log_2 C_k \sim k^2$; the same leading entropy persists after reduction modulo any fixed $M \geq 2$. A corresponding affine bound changes the pattern loss in a localized phase-removal lemma from $2^{-O(k^3)}$ to $2^{-2k^2 - O(k \log k)}$. For progression counts, a Cauchy–Schwarz induction retaining two distinguished functions gives quadratic, rather than linear, dependence on the balanced uniformity norm, with explicit density factors. We also solve a weighted finite-partition density-extraction problem exactly, including its fixed-cell-count extremizers, and give a density-dependent iteration potential with an asymptotically optimal leading cost for a specified abstract recurrence. Complete proofs, computational certificates,

a source-level Lean interface map, and a further-research agenda are included. These results strengthen the indicated auxiliary statements; a new global bound for Szemerédi’s theorem is not established.

The most productive way to improve a long quantitative proof is to isolate the places where information is discarded, determine which losses are necessary, and then propagate the improved statements through the actual dependencies. This report carries out that program at four points of Gowers’ argument [1]: finite floor-pattern counting, uniformity-based progression counting, selection of a dense partition cell, and the final density iteration.

The strongest combinatorial conclusion concerns the maps

$$\varepsilon \mapsto \left\lfloor \sum_{i=1}^k \varepsilon_i \alpha_i \right\rfloor, \quad \varepsilon \in \{0, 1\}^k, \quad -r < \alpha_i < r.$$

Their number is exactly polynomial in the integer slope range r : it is $(2r)^k C_k$. The remaining coefficient satisfies $\log_2 C_k \sim k^2$. Thus the range dependence is determined exactly, and the leading quadratic exponent in the dimension is determined as well. The lower bound uses the known threshold-function asymptotic of Zuev [9]; we give the reduction in full. This is a concrete resolution of the counting subproblem used here, rather than a claim to have settled a major open problem about arithmetic progressions.

The analytical refinement is equally explicit. If A has density δ in a suitable finite abelian group and $\epsilon = \|\mathbf{1}_A - \delta\|_{U^{k-1}}$, then

$$\left| \Lambda_k(\mathbf{1}_A) - \delta^k \right| \leq \sqrt{\delta} \epsilon^2 \sum_{t=0}^{k-3} (t+1) \delta^t \quad (k \geq 3). \quad (\text{F.11.1})$$

The group need not have prime order: invertibility of the relevant slope differences suffices. For three-term progressions the factor $\sqrt{\delta}$ can be replaced by $\sqrt{\delta(1-\delta)}$, and singletons give equality. The Cauchy–Schwarz and telescoping mechanisms are standard; the contribution here is their explicit, stronger formulation and its careful comparison with the displayed source bounds.

F.11.1 Which source is being strengthened?

We inspected both user-specified paths in the ProveIt repository and pinned the discussion to commit 7d99e5deeb209a4671060990cb1ea3d672d57064 [41]. Its `sz-thm-gowers-proof.tex` is an edited transcription with mathematical and rounding corrections. A reference below to a numbered “source” statement means that inspected version unless the journal paper is specified separately. This distinction matters: the edited source explicitly records that a changed estimate in Corollary 16.11 leaves the displayed quantitative proof of Theorem 18.1 unsupported. That is a statement about this derivation and formalization checkpoint, not a refutation of Gowers’ published theorem.

The repository’s checkpoint reports 95 exact companion proofs among 120 numbered statements, with 25 still open. It also reports that 140 reachable modules compiled and that its axiom audit checked 1,365 public theorems. These are source-reported results, not a fresh Lean execution in this report. In particular, a proposition-valued catalogue definition and a proved conditional implication have different mathematical statuses from an unconditional proof of the catalogue entry. Section V.26 records the interfaces relevant to the present work.

[write] 6 October 2026; sources 06–10. *Stale since the pin.* These figures are those of the pin 7d99e5dee (its `FORMALIZATION_STATUS.txt` and ledger). At HEAD 764740f07 (6 October 2026) the ledger records 99

exact companions and 21 open statements, and `FORMALIZATION_STATUS.txt` reports 196 reachable modules and 1,805 audited public theorems. The 4 statements closed since the pin are `lemma_16_1`, `lemma_16_5`, `lemma_16_6` and `lemma_16_9`. `corollary_16_11`, `theorem_18_1` and `theorem_18_2` are still open, so the dependency warning above stands. Crosswalk: Section V.1.

F.11.2 Results and their precise status

Source location	Refinement proved here	What it establishes
Corollaries 17.5–17.6	Exact integer/fractional decomposition; sharp leading carry entropy	Polynomial range dependence, a quadratic dimension exponent, and modular saturation
Proposition 17.7	Carry-count bound propagated through localized Cauchy–Schwarz	An explicit improved local phase-removal loss
Theorem 3.2 and Corollaries 3.3, 3.6	Two distinguished norms, double telescoping, exact diagonal subtraction	Stronger modular counting and existence criteria
Lemma 5.15	Weighted asymmetric extraction with extremizers	Optimal mass and fixed-cell-count guarantees from the stated data
Section 18 iteration	Explicit discrete potential	An optimal leading cost for a stated abstract local recurrence

The results are proved below and are quantitative strengthenings relative to those source formulations. We make no historical priority claim for every inequality or decomposition. A bounded literature search cannot establish such a claim. The comparison with modern work is also important: Leng, Sah, and Sawhney proved, for $k \geq 5$, a bound of the shape

$$r_k(N) \ll_k N \exp(-(\log \log N)^{c_k}), \quad c_k > 0,$$

using quasipolynomial inverse-theorem bounds [2, 3]. Our auxiliary refinements do not improve that result. They are useful as standalone estimates, as replacements in a carefully tracked proof, and as tractable formalization targets.

F.11.3 Conventions

All averages on finite sets are normalized unless a sum is written explicitly. Logarithms are natural except for $\log_2$. We write

$$\mathcal{R}_d(m) = \sum_{j=0}^d \binom{m}{j},$$

where $\binom{m}{j} = 0$ for $j > m$. The standard hyperplane-arrangement bound is that m affine hyperplanes in $\mathbb{R}^d$ have at most $\mathcal{R}_d(m)$ complementary regions. Adding one hyperplane and counting the regions it cuts proves the recurrence $R(m, d) \leq R(m-1, d) + R(m-1, d-1)$, and induction gives the binomial sum. When $m \geq d \geq 1$,

$$\mathcal{R}_d(m) \leq (em/d)^d. \tag{F.11.2}$$

For completeness, take $z = d/m \leq 1$. Since $z^j \geq z^d$ for $j \leq d$, the binomial theorem gives $z^d \mathcal{R}_d(m) \leq (1+z)^m \leq e^d$.

Floor patterns on hyperplanes require more care than patterns in open regions. A one-sided perturbation lemma below supplies the needed justification; counting only open regions without that step would leave a genuine gap.

F.12 Source 07. The problem and the concrete improvements

Source 07, delivered Section 1.

Added in the second write, 6 October 2026. Delivered title block. “Additive combinatorics / Quantitative proof refinement / Transport-sensitive rigidity in Gowers’s approximate-homomorphism argument / Higher-order Freiman structure without square-root losses / Research draft / 6 October 2026”; PDF author field “Research draft”. Archive `Gowers_Transport_Rigidity.zip`, manuscript `transport_rigidity.tex` (1,472 lines, 21 pages), pin `ae0e89574`. Printed below: the abstract, the title-page statements and Section 1; Appendix A and Sections 2–8 form a chapter of Part II, Sections 9–11 and Appendix B are in Part V.

Abstract of source 07 (as delivered). We replace a nested almost-everywhere argument by a coupled-path argument that keeps local errors in their averaged form. The resulting theorem applies to finite fibered domains, nonuniform fiber sizes, missing fibers, and arbitrary abelian target groups. Its exact error budget is expressed by a finite transport profile; coarse versions require $2[1 + (k - 1)\kappa]\delta < 1$, or $[2k + \tau k(k - 1)]\delta < 1$ under additive fiber-size invariance, to force a Freiman homomorphism of order k . For constant fiber sizes the strict threshold $\delta < 1/(2k)$ is optimal, even with identical cyclic source and target groups.

Applied to the hypotheses recorded for Gowers’s Lemma 10.9, the argument upgrades order two to order five without a stronger smallness assumption. Retaining the ambient fiber-invariance assumptions of the original construction upgrades it further to order nine. We also give explicit local relaxations of the error allowance, a collision-to-mode criterion, a nonabelian word-relation extension, and a rational linear program for optimizing the starting law. Every asserted general result is proved in the text. Exact finite tests are included as supplementary checks, not as replacements for the proofs.

Scope. This is a local strengthening of an exactification step, not a new upper bound for progression-free sets. The sharp threshold problem solved here is explicitly formulated in this note. Priority over the full literature has not been established. The new results have not been compiled in Lean or independently refereed.

Repository baseline. The comparison is pinned to `ProveIt` commit `ae0e895748209f0dd079c6b5ca6df6c6fc065a81`; it does not depend on the subsequent state of `main`.

Package. LaTeX source, compiled PDF, exact-arithmetic Python checks, their output, a source manifest, and a formalization plan.

A recurring task in additive combinatorics is to convert a nearly consistent system of local differences into an exactly additive map. Suppose that a candidate value $\psi(d)$ usually equals $\phi(y) - \phi(x)$ when the index of y is the index of x plus d . A relation

$$a + b = c + d$$

can then be tested by comparing two paths with the same initial and terminal indices. If all four edges are correct, their sums agree. The substantive issues are which probability law to put on the paths, how that law changes under translation, and how to treat a path that leaves the usable domain.

Gowers’s Lemma 10.9 in [1] obtains an order-two Freiman homomorphism from the local model of Lemma 10.6 under

$$\theta = 10\eta^{1/5}, \quad 6\sqrt{\theta} < 1. \tag{F.12.1}$$

The repository records the relevant assumptions explicitly and contains a companion proof [42, 43]. We work at that precise interface. The replacement keeps a single edge-defect budget rather than successively balancing outer and inner exceptional sets.

F.12.1 What is proved

Write $\text{Fr}_k(D, H)$ for the maps from D to H preserving every equality between two sums of k elements of D , repetitions allowed. The main results are as follows; all parameters are defined in detail below.

Input or setting	Conclusion	Location
Nonuniform finite fibers	Every violated word relation has transported error budget at least one.	Theorem II.15.1
Fiber ratio at most κ	$2[1 + (k - 1)\kappa]\delta < 1$ forces Fr_k .	Corollary II.15.3
Constant fiber sizes	The threshold $\delta < 1/(2k)$ is optimal for arbitrary direction sets, for every $k \geq 2$.	Theorem II.16.1
Recorded Lemma 10.9 hypotheses	Order five, rather than merely order two.	Corollary II.17.3
Full ambient fiber invariance	Order nine under the same original parameter range; $\eta \leq (40k)^{-5}$ suffices for order k .	Theorem II.17.5, Corollary II.17.6
Local order-two exactification	Explicit allowances 120^{-5} and 75^{-5} , under two different hypothesis packages.	Propositions II.17.4, II.17.7

The numbers five and nine are certified consequences, not claims that these are the largest possible orders under the respective assumptions. The sharp claim concerns the uniform directional defect δ , not the original parameter η , which has already passed through an earlier construction.

F.12.2 Position relative to other results

Homomorphism testing and correction have an established literature. In particular, the majority-difference method and optimal average-test thresholds are treated in [27, 28]. Our assumptions are different: we control each selected direction in a possibly punctured, fibered domain, and keep explicit track of the law of the source point. Thus the threshold proved below should not be read as a replacement for a sharp *average* homomorphism test.

There is also a substantial distinction between refining one lemma and improving a final density bound. For context, the 2024 theorem of Leng, Sah, and Sawhney gives, for each fixed $\ell \geq 5$, an estimate of the form $r_\ell(N) \ll N \exp[-(\log \log N)^{c_\ell}]$, with $c_\ell > 0$ [2]. We do not improve that estimate or assert that the local gains below survive the entire quantitative iteration. No claim about a most recent record bound is needed for the arguments in this note.

F.13 Source 08. The target lemma and the contribution

Source 08, delivered Section 1.

Added in the second write, 6 October 2026. Delivered title block. “Additive combinatorics / Research memorandum / Fourth-order stability of additive cube counts / Torsion-sensitive refinements of Gowers’ Lemmas 3.7 and 3.10 / Prepared for Vladimir Reshetnikov / 6 October 2026”; PDF author field “Research memorandum prepared for Vladimir Reshetnikov”. Archive `gowers_cube_refinements.zip`, manuscript `fourth_order_cube_counts.tex` (1,677 lines, 25 pages). It records no commit: its `source_provenance.json` gives blob identifiers, which are those of `7d99e5dee` (checked at intake), and its links follow `main`. Printed below: the abstract, the title-page status and Section 1; Sections 2–10 and Appendix A are in Part III, Sections 11–14 and Appendix B in Part V.

Abstract of source 08 (as delivered). For a finite abelian group G , write $p = 2^d$, $d \geq 2$, and let $\mathcal{C}_d(F)$ be the normalized count of d -dimensional additive cubes weighted by a real function $F = \delta + f$, where $\mathbb{E}f = 0$. The first three nonconstant terms of this cube-count polynomial vanish identically. Consequently, with $t = \|f\|_{U^d(G)}$, we obtain

$$0 \leq \mathcal{C}_d(F) - \delta^p \leq \sum_{r=4}^p \binom{p}{r} \delta^{p-r} t^r.$$

This removes the linear, quadratic, and cubic losses in the Minkowski bound used in Gowers' Lemma 3.10. In the paper's uniformity parameter $\alpha = t^p$, the resulting exponent is $4/p$, and this exponent cannot be increased uniformly, even for genuine subsets of cyclic groups of odd prime order.

We determine the exact fourth-order coefficient over every finite abelian group. It is $R_d \|f\|_{U^2}^4 + T_d E_{\langle 2 \rangle}(f)$, where R_d counts ordinary parallelograms in the Boolean cube, T_d counts additional binary affine planes, and $E_{\langle 2 \rangle}(f)$ is the fourth Fourier moment on nontrivial order-two characters. We also give the complete three-cube expansion over odd-order groups, identify its signed fifth-order term, and construct a counterexample to a natural global lower bound suggested by the fourth-order coefficient. Further results include a Fourier stability inequality, an inhomogeneous version, and a linear-algebraic criterion for which centered terms can survive.

Status. The main statements below have mathematical proofs; finite identities are additionally checked by the accompanying exact-arithmetic program. These are local refinements of cube-count estimates, not a new global bound for Szemerédi's theorem. This AI-assisted research draft has not undergone independent peer review. Bibliographic priority and Lean verification are not asserted.

F.13.1 What is being strengthened

Gowers' *A new proof of Szemerédi's theorem* [1] contains two complementary estimates for additive cubes. Lemma 3.7 gives a lower count $\delta^{2^d} |G|^{d+1}$. In the cyclic setting of the paper, Lemma 3.10 gives the upper count

$$(\delta + \alpha^{1/2^d})^{2^d} |G|^{d+1} \tag{F.13.1}$$

when the set is α -uniform of degree $d - 1$. The degree shift is important: a d -dimensional cube corresponds to the normalized U^d norm, while the paper calls the relevant uniformity *degree* $d - 1$. The displayed bound follows from Minkowski's inequality.

The supplied Lean file `Proofs03CubeUpper.lean` proves this statement directly with the unnormalized norm, applying `gowersNorm_add_1e` [46]. We keep its conclusion as our baseline rather than replacing it with a looser, already expanded estimate. The improvement is that centering eliminates the first three terms that Minkowski would retain.

Throughout the article, α may be taken to be the actual normalized balanced cube moment $\|f\|_{U^d}^{2^d}$. If the hypothesis supplies only an upper bound on that moment, every upper estimate remains valid after replacing t by $\alpha^{1/2^d}$.

F.13.2 A guide to the results

The central results are organized by the kind of information they use. Theorem III.7.3 uses only centering and Gowers–Cauchy–Schwarz. Theorem III.8.3 additionally uses the arithmetic of four

cube vertices, and distinguishes odd groups from groups with order-two characters. Theorem III.9.2 computes every coefficient of the three-cube polynomial in terms of a small number of configuration averages. Theorem III.11.4 proves that the improved exponent is already optimal in the original indicator-function setting.

Two conclusions are worth keeping separate. The exponent $4/2^d$ is optimal for a uniform estimate in the balanced U^d moment. The exact leading coefficient is a Fourier quantity, not merely a universal multiple of $\|f\|_{U^d}^4$. In particular, the coefficient obtained by replacing the Fourier quantity with $\|f\|_{U^d}^4$ need not be sharp on odd-order groups.

F.13.3 Relation to existing work and limits of the claim

The method is an explicit use of familiar tools: multilinear expansion, character orthogonality, and repeated Cauchy–Schwarz. The broader study of linear configurations includes the true-complexity program of Gowers and Wolf [29], Manners’ polynomial bounds obtained through iterated Cauchy–Schwarz [14], and the study of norming linear systems by Cho, Conlon, Lee, Skokan, and Versteegen [30]. We do not claim to introduce any of these general methods.

The specific contributions here are the sharpened statement relative to (F.13.1), its optimal exponent for sets, the all-dimension quartic formula with its order-two correction, and the explicit signed three-cube analysis. A targeted literature check did not establish priority for this package of formulas. They should therefore be described as *refinements derived in this memorandum*, not as certified first results.

For comparison, Leng, Sah, and Sawhney prove, for $k \geq 5$, a bound of the form $r_k(N) \ll_k N \exp(-(\log \log N)^{c_k})$ with $c_k > 0$ [2]. Nothing below improves that global result. A stronger cube-count estimate does not by itself improve the inverse theorem or the density-increment iteration required to count long arithmetic progressions.

F.14 Source 09. Scope, source theorem, and principal results

Source 09, delivered Section 1.

Added in the second write, 6 October 2026. Delivered title block. “Arithmetic Circuit Expansions for Cubical Counts / A fourth-order strengthening of Gowers’ Lemma 3.10, exact torsion corrections, and a formalization roadmap / Research manuscript prepared with ChatGPT / 6 October 2026”; the PDF author field reads “Research manuscript prepared with ChatGPT”. Archive `gowers_cubical_expansions.zip`, manuscript `article.tex` (1,794 lines, 27 pages), pin `7d99e5dee` (its source audit, shipped as `09-cubical-expansions-SOURCE_AUDIT.md`). Printed below: the abstract and Section 1; Sections 2–10, 13 and Appendices A–B are in Part III, Sections 11, 12, 14, 15 and Appendix C in Part V. Its ‘question’ environment, printed “Research question” there, is set in this report’s ‘Research question’ style.

Abstract of source 09 (as delivered). Let G be a finite abelian group, $d \geq 2$, $m = 2^d$, and $F = \delta + f$ a real function with $\mathbb{E}f = 0$ and $\varepsilon = \|f\|_{U^d}$. The triangle-inequality estimate behind Gowers’ Lemma 3.10 bounds the normalized d -cube count by $(\delta + \varepsilon)^m$. We retain the balancing condition and prove that all terms of degrees one, two, and three vanish. Consequently the error is fourth order, with a sharp exponent, rather than first order. We then classify every four- and five-vertex contribution over arbitrary finite abelian groups. The exact fourth-order term is $\delta^{m-4}(P_d E_4 + T_d T_2)$, and the fifth-order term is $\delta^{m-5}(Q_d M_5 + S_d T_{3,5})$; the four coefficients are explicit Stirling-number expressions. The additional statistics T_2 and $T_{3,5}$ are supported on characters of orders two and three. The remaining error is bounded by an explicit sixth-order

binomial tail. We prove sharpness statements, construct a negative fifth-order correction, transfer weighted examples to actual sets, and exhibit a six-vertex determinant-5 configuration showing the first possible order-five torsion contribution. Exact finite computations and a repository-specific formalization plan accompany the proofs.

These are local counting results proved in this manuscript. Historical originality has not been established, and the manuscript has not been refereed or verified by a proof assistant. No improved final bound for Szemerédi’s theorem is claimed.

F.14.1 The precise target

The target is Lemma 3.10 of Gowers’ *A new proof of Szemerédi’s theorem* [1]. In the notation of that paper, a set of density δ which is α -uniform of degree $d - 1$ has at most

$$(\delta + \alpha^{1/2^d})^{2^d} N^{d+1} \tag{F.14.1}$$

parameterized d -dimensional cubes. The proof applies the norm triangle inequality to the indicator written as its mean plus its balanced part. The present paper asks what is lost when the mean-zero property is forgotten at that step.

The supplied repository contains both a transcription of the paper and a Lean development [47]. At the inspected commit, `Proofs03CubeUpper.lean` proves `lemma_3_10_holds` using exactly this norm decomposition. Our comparison is with that theorem, not with an imagined gap in its proof. The repository’s definitions use unnormalized Fourier transforms and unnormalized Gowers norms. Section V.38 gives the precise conversion and the inspected source identifiers.

F.14.2 The first strengthening

Write $m = 2^d$, let $F : G \rightarrow [0, 1]$, and put $\delta = \mathbb{E}F$, $f = F - \delta$, and $\varepsilon = \|f\|_{U^d(G)}$. We prove

$$0 \leq C_d(F) - \delta^m \leq \sum_{s=4}^m \binom{m}{s} \delta^{m-s} \varepsilon^s. \tag{F.14.2}$$

Thus, at fixed positive density and dimension, the error in (F.14.1) improves from $O(\alpha^{1/m})$ to $O(\alpha^{4/m})$. This comparison concerns the excess above δ^m ; it does not mean that all constants are uniform as d increases or δ decreases. The full polynomial in (F.14.2) is the appropriate statement in those regimes.

The exponent four cannot be increased, even for actual sets in growing finite groups. Moreover, the coefficient of the fourth-order term can be substantially smaller than $\binom{m}{4}$. Set

$$A_d = \frac{m(m-1)(m-2)}{24}. \tag{F.14.3}$$

The exact classification implies a universal upper bound with $A_d \delta^{m-4} \varepsilon^4$ as its fourth-order term. This coefficient is sharp when all finite abelian groups are permitted. Compared with $\binom{m}{4}$, it saves the factor $m - 3$.

F.14.3 The complete expansion through degree five

For $d \geq 3$ we prove

$$\begin{aligned} \mathcal{C}_d(\delta + f) &= \delta^m + \delta^{m-4}(P_d E_4(f) + T_d T_2(f)) \\ &\quad + \delta^{m-5}(Q_d M_5(f) + S_d T_{3,5}(f)) + R_6, \end{aligned} \quad (\text{F.14.4})$$

where the Fourier statistics are defined in Section III.21, and

$$\begin{aligned} P_d &= 2^{d-2} \begin{Bmatrix} d+1 \\ 3 \end{Bmatrix}, & T_d &= 2^{d-2} \begin{Bmatrix} d+1 \\ 4 \end{Bmatrix}, \\ Q_d &= 2^d \begin{Bmatrix} d+1 \\ 4 \end{Bmatrix}, & S_d &= 2^d \begin{Bmatrix} d+1 \\ 5 \end{Bmatrix}. \end{aligned} \quad (\text{F.14.5})$$

Here $\begin{Bmatrix} n \\ k \end{Bmatrix}$ is a Stirling number of the second kind. The remainder satisfies

$$|R_6| \leq \sum_{s=6}^m \binom{m}{s} \delta^{m-s} \varepsilon^s \leq \binom{m}{6} \varepsilon^6 (\delta + \varepsilon)^{m-6}. \quad (\text{F.14.6})$$

All identities are exact finite identities, before any limiting process. The formula does not require a prime modulus, distinct evaluated cube vertices, or a lower bound on $|G|$.

There are two genuinely different sources of terms. Real affine relations among cube vertices give E_4 and M_5 . Integer lattice indices give the torsion terms T_2 and $T_{3,5}$. Ignoring this arithmetic distinction would give false formulas on groups with small prime torsion.

F.14.4 Boundaries of the claims

The proofs below establish the stated identities, estimates, examples, and finite-configuration classifications. We do not assert that these identities are absent from all previous literature. The broad framework of multilinear Gowers inequalities and norm geometry is established work; see, for example, Gowers [1] and Hatami [31]. The contribution pursued here is a specific, fully explicit local refinement, including its arithmetic corrections and sharpness analysis.

An improved auxiliary cube estimate is not automatically an improved arithmetic-progression theorem. A new global bound would require a complete quantitative argument connecting this refinement to an inverse step and a density-increment iteration. No such argument is supplied. Current results and the status of the paper's conjectural progression statements are discussed separately in Section III.25.

F.15 Source 10. Scope, results, and the baseline

Source 10, delivered Section 1.

Added in the second write, 6 October 2026. Delivered title block. “Research report / Quantitative Refinements Around Gowers’ Proof of Szemerédi’s Theorem / Torsion, density profiles, additive energy, and sharp limits for polynomial colorings / A mathematical research report / 6 October 2026”; PDF author field “Research report”. Archive `gowers_refinements.zip` (independent of source 06’s `Gowers_Refinements(1).zip`; its manuscript is not the file of the same name of sources 01 and 02), manuscript `gowers_refinements.tex` (2,443 lines, 34 pages), pin 7d99e5dee. Printed below: the abstract, the title-page note and Section 1; Sections 2–4 are in Part I, 7–8 in Part II, 5, 6 and 9 in Part III, 10–12 and Appendix A in Part V.

Abstract of source 10 (as delivered). We develop explicit strengthenings of several finite steps surrounding Gowers’ 2001 proof of Szemerédi’s theorem, using the supplied paper and Lean development as a precisely identified baseline. The results include an optimal conversion from partition discrepancy to a density increment, a density-sensitive Fourier-to-progression tradeoff with all integer rounding explicit, a mixed- L^p generalized von Neumann inequality on finite abelian groups with quantified torsion losses, and a density-weighted counting estimate replacing an exponential expansion loss. A weighted energy argument and a sharper vertical-fibre extraction improve the explicit order-eight Freiman restriction from $2^{-1882}\gamma^{1164}$ to $2^{-209}\gamma^{103}$ of the original domain. We also determine exactly the feasible dimensions for a field-coordinate polynomial coloring that excludes symmetrically colored even k -term progressions, when every coordinate degree is at most k and the prime characteristic exceeds k . Complete proofs, sharp examples, finite diagnostic checks, a formalization map, and a proposed research program are provided. These are local refinements and a restricted extremal theorem; the report does not claim a new best general bound for Szemerédi’s theorem or priority for methods already present in the literature.

Source snapshot: 7d99e5deeb209a4671060990cb1ea3d672d57064. All newly stated results are proved in this report. The accompanying computations are diagnostic supplements, not Lean kernel certification.

The most productive response to a request to strengthen a long proof is to identify exactly which hypotheses can be weakened, which losses can be reduced, and which parts of the argument obstruct further progress. This report carries out that program for several stages of Gowers’ proof [1]. Its main outputs are explicit finite theorems, with enough detail to support later insertion into the accompanying formal development.

F.15.1 The source version matters

The two supplied repository paths were read at commit

7d99e5deeb209a4671060990cb1ea3d672d57064.

Its recorded commit time is 6 October 2026, 19:47:40 UTC [48]. The repository’s $\text{\LaTeX}$ is an edited transcription: it includes editorial corrections and qualifications and is not identical to the published 2001 article. For example, the proof of Corollary 2.5 already contains a stronger intermediate density and length estimate than its displayed conclusion. Theorem 3.2 has explicit prime-modulus and range assumptions addressing a noninvertible reindexing step. We compare against these existing repairs where relevant.

The supplied status ledger records 120 numbered catalogue statements: 95 have exact companion proofs and 25 remain open. Its axiom-audit source explicitly distinguishes an axiom check from a check that the hypotheses of conditional results have been discharged. We have inspected these source records; we have not rerun the Lean project or certified the new theorems in Lean. Nothing below treats an unproved catalogue declaration as an established mathematical input.

[write] **6 October 2026; sources 06–10.** *Stale since the pin.* The counts 95 and 25 are those of the pin 7d99e5dee. At HEAD 764740f07 (6 October 2026) the ledger records 99 exact companions and 21 open statements; the 4 statements closed since the pin are `lemma_16_1`, `lemma_16_5`, `lemma_16_6` and `lemma_16_9`. Crosswalk: Section V.1.

F.15.2 A comparison of the main finite guarantees

In this table δ is density, η a normalized correlation or Gowers norm as specified in the relevant theorem, c an energy proportion, and γ a graph-energy proportion. The meanings are not interchangeable.

Step	Baseline in the supplied argument	Refinement proved here
Partition selection	Increment at least $S/2$ from total balanced discrepancy S	The optimal increment $\delta S/(2\delta - S)$; an exact tail bound and an exceptional-set version (Section I.23).
Fourier extraction, Corollary 2.5	The proof already gives increment $\eta/4$ and length $\sqrt{\eta N/(128\pi)}$ in its large-scale range	The density-sensitive formula (I.24.1), increment $\delta\eta/(4\delta - \eta)$, and a complete tradeoff with no size side condition (Section I.24).
Generalized von Neumann, Theorem 3.2	Prime modulus, $k \leq N$, and 1-bounded functions	Arbitrary finite abelian groups, explicit kernel factors, general coefficient systems with an invertible target-neighbor difference, and mixed L^p norms (Theorem III.28.2).
Progression counting, Corollary 3.3	Error bounded by $2^k\alpha$ after normalization	A weighted telescoping sum, at most $(k-2)\alpha$ for $k \geq 3$ in the same no-torsion-loss setting, with additional density factors (Section III.29).
Energy extraction, Proposition 7.3	Retain $2^{-20}c^{12}m$; difference set at most $2^{38}c^{-24}m$	Retain $cm/4$; difference set at most $2^{10}c^{-5}m$ (Theorem II.20.1).
Freiman extraction, Lemma 7.5	Retain $ B /(8kC^{4k})$	Retain $ B /[k(h_k - 1) + 1]$, and hence $ B /(kC^{2k+1})$ (Theorem II.21.1).
Order eight, Corollary 7.6	Retain $2^{-1882}\gamma^{1164} B_0 $	Retain $2^{-209}\gamma^{103} B_0 $ (Corollary II.21.2).
Polynomial coloring restriction	A layered-norm construction and a qualitative Chevalley–Warning limitation in recent work	An exact dimension criterion for prescribed coordinate degrees $1 \leq D_\ell \leq k$ over $\mathbb{F}_p$, $p > k$ (Section III.30).

F.15.3 What sort of contribution these results make

The partition cap and tail bounds are sharp finite optimization statements. The torsion factor is exactly sharp for three-term progressions in elementary abelian 2-groups and has the correct group-size order for a family of longer patterns. The polynomial dimension theorem gives both a universal obstruction in a precisely specified class and a construction attaining it. These are explicit mathematical conclusions of the proofs below.

The underlying tools—Cauchy–Schwarz, Fourier orthogonality, convexity, dependent random choice, sumset inequalities, and Chevalley–Warning—are established. The matching norm-layer coloring is credited to the recent literature. The exact formulations assembled here may be useful new entries in this project, but no assertion of publication-level novelty is made without a broader priority investigation. In particular, an improvement over a 2001 constant is not automatically an improvement over the strongest known theorem in 2026.

F.15.4 Relevant contemporary benchmarks

Let $r_k(N)$ be the largest size of a subset of $\{1, \dots, N\}$ without a nonconstant k -term progression. Gowers’ historical general bound has the shape $N(\log \log N)^{-c_k}$ [1]. Relevant later benchmarks, checked against primary sources for this report, include

$$\begin{aligned} r_3(N) &\leq N \exp\left(-c \frac{(\log N)^{1/6}}{\log \log N}\right) && \text{(Raghavan, 2026 preprint),} \\ r_4(N) &\ll N(\log N)^{-c} && \text{(Green–Tao),} \\ r_k(N) &\ll_k N \exp(-(\log \log N)^{c_k}), \quad k \geq 5 && \text{(Leng–Sah–Sawhney).} \end{aligned}$$

See [34, 6, 2] for the precise theorems and their hypotheses. This is a set of comparison benchmarks, not a claim that the present report reproduces their proofs.

The status of the old uniform-set conjectures must also be updated. Gowers’ later construction disproves the random four-term lower bound under Fourier uniformity [35]. The 2026 preprint of Shi and Dong reports counterexamples to the corresponding conjectured lower bound for every even length at least four [33]. Those statements should not be proposed again as untouched open conjectures. Their coloring construction and the unresolved questions around it instead motivate the restricted extremal problem in Section III.30.

F.16 Source 11. The problem and the principal results

Source 11, delivered Section 1.

Added in the third write, 6 October 2026. Delivered title block. “Sharp Affine Rigidity from Graph Additive Energy / Cubic stability, target torsion, and exact non-affine extremizers / A quantitative refinement of the Section 6 interface in Gowers’s proof / Research article prepared for Vladimir Reshetnikov / 6 October 2026”; the PDF author field reads “Research article prepared for Vladimir Reshetnikov” and the running head “Research note”. Archive `Gowers_Affine_Rigidity_Research.zip` (arrival 08ab4187e), manuscript `article.tex` (1,259 lines, 19 pages), pin 40f885098 (its `sources.json`, shipped as `data/11-affine-rigidity-sources.json`). Printed below: the abstract and Section 1; Sections 2–11 form a chapter of Part II (Section II.22), Sections 12–14 and Appendix A are in Part V. Its ‘question’ environment, printed “Research question” there, is set in this report’s ‘Research question’ style.

Abstract of source 11 (as delivered). The exact-additivity observation adjacent to Lemma 6.2 of Gowers’s *A new proof of Szemerédi’s theorem* says that a function with the maximum possible number of respected additive quadruples is affine. We develop a sharp quantitative version near this endpoint. For a map $f : G \rightarrow W$ between finite abelian groups, put $N = |G|$ and let $\mathcal{E}(f)$ count ordered additive quadruples respected by f . If f differs from an affine map on a proportion $t < 1/7$ of its domain and W has no nonzero element of order two, we prove

$$1 - \frac{\mathcal{E}(f)}{N^3} \geq 4t - 10t^2 + 6t^3.$$

Equality holds precisely when the error has one nonzero value on a coset of a subgroup. We prove a stronger inequality separating the energy deficit of the error support from the dispersion of its nonzero values. If the target contains an involution, the sharp cubic is instead $4t - 12t^2 + 8t^3$, with involution-valued errors in the equality case. An elementary correction argument removes the initial proximity assumption when the graph-energy defect is less than $1/9$. Consequently, for arbitrary finite abelian domains of size at least 34, the exact largest non-affine graph energy and all its extremizers are determined. A second proof removes this size restriction for maps

between vector spaces over the same prime field, apart from domains on which every map is affine. The paper includes a partial-domain extension, a high-correlation corollary to Gowers's Proposition 6.1, all-orders inversion of the odd-target cubic, finite verification, and a repository-specific formalization plan. These are local structural refinements; no improved global bound in Szemerédi's theorem or completed Lean formalization is claimed.

F.16.1 Which part of Gowers's argument is being refined?

In Section 6 of Gowers's paper, Fourier information about first differences is converted into additive energy of the graph of a frequency-selection map. Proposition 6.1 supplies the quantitative energy lower bound; Lemma 6.2 identifies the respected quadruples with additive quadruples in the graph. The observation immediately preceding that lemma treats the exact endpoint: full energy forces a full domain and an affine map [1, pp. 498–501]. The question here is how much of this rigidity survives just below full energy, with sharp constants and equality cases.

The graph-counting identity itself cannot be made stronger: it is an identity. Our contribution is a quantitative structural conclusion at the same interface. In a sufficiently high-energy regime, one can recover a single affine map on almost the entire domain rather than merely seeking a structured piece through the subsequent Freiman-type arguments. The hypothesis is much stronger than the small-energy hypotheses needed in the general proof, a distinction that will remain explicit.

F.16.2 Notation for the main statements

Throughout, G and W are finite abelian groups, written additively, and W is nontrivial. Set $N = |G|$. An affine map is a map $\ell(x) = L(x) + b$, with $L \in \text{Hom}(G, W)$ and $b \in W$. For $f : G \rightarrow W$ define

$$\mathcal{E}(f) = \#\{(x_1, x_2, x_3, x_4) \in G^4 : x_1 + x_2 = x_3 + x_4, \\ f(x_1) + f(x_2) = f(x_3) + f(x_4)\}, \quad (\text{F.16.1})$$

$$\varepsilon(f) = 1 - \mathcal{E}(f)/N^3, \quad (\text{F.16.2})$$

$$d(f) = \min_{\ell \in \text{Aff}(G, W)} \frac{|\{x : f(x) \neq \ell(x)\}|}{N}. \quad (\text{F.16.3})$$

All quadruples are ordered and repetitions are allowed. In particular, there are exactly N^3 additive quadruples in G .

Let $\kappa = 0$ when W has no nonzero element of order two, and let $\kappa = 1$ otherwise. Define

$$\Psi_\kappa(t) = 4t - (10 + 2\kappa)t^2 + (6 + 2\kappa)t^3. \quad (\text{F.16.4})$$

Both cubics are strictly increasing on $[0, 1/7]$. We write D_κ for the inverse of this restricted branch. Its domain includes $[0, 1/9]$.

Theorem F.16.1 (Sharp small-error profile). *Suppose that f differs from an affine map ℓ on a set S of density t .*

(i) *If $\kappa = 0$ and $0 < t < 1/7$, then $\varepsilon(f) \geq \Psi_0(t)$. Equality holds if and only if*

$$f = \ell + b\mathbf{1}_{a+H}$$

for a subgroup $H \leq G$, an element $a \in G$, and a nonzero $b \in W$, where $S = a + H$.

(ii) If $\kappa = 1$ and $0 < t < 1/3$, then $\varepsilon(f) \geq \Psi_1(t)$. Equality holds if and only if the same description holds with b a nonzero element of order two.

The assertions for $t = 0$ reduce to $\mathcal{E}(f) = N^3$.

[write] **7 October 2026; sources 54–59.** The first range extended by source 59 (Section II.248): for targets without involutions $\varepsilon(f) \geq \Psi_0(t)$ holds on $0 < t < 4/9$, and on $0 < t < 1/2$ without 3-torsion as well, with the same coset equality cases (its Theorem II.250.4 and Corollary II.250.6); the colour coefficient $2(1 - 7t)$ of Theorem II.25.3 becomes $2(1 - 3t)$ (its Theorem II.249.3). The involution case (ii) is not extended.

The two sufficient radii, $1/7$ and $1/3$, are not asserted to be optimal. The cubic coefficients and the stated equality families are exact. In particular, “sharp” does not mean that every real t occurs as the density of a subgroup coset in every fixed group.

Theorem F.16.2 (Unconditional high-energy correction). *If $\varepsilon(f) < 1/9$, then*

$$d(f) \leq D_\kappa(\varepsilon(f)). \quad (\text{F.16.5})$$

The closest affine map is unique. If $d(f) > 0$, equality in (F.16.5) is equivalent to the applicable subgroup-coset error description in Theorem F.16.1.

[write] **7 October 2026; sources 54–59.** For targets without involutions source 59 extends the range to $\varepsilon(f) < 2/9$, combining Theorem F.16.1 with the published correction theorem of Ben-Or, Coppersmith, Luby and Rubinfeld, as it credits (its Theorem II.251.2).

For odd-order targets, the small-defect expansion is

$$D_0(u) = \frac{u}{4} + \frac{5u^2}{32} + \frac{11u^3}{64} + \frac{475u^4}{2048} + O(u^5). \quad (\text{F.16.6})$$

Thus the leading recovery constant is $1/4$, not merely an unspecified absolute constant. Each displayed coefficient is forced along families of coset errors whose densities tend to zero.

Theorem F.16.3 (Exact largest non-affine energy). *For any finite abelian G, W as above with $N \geq 34$, every non-affine $f : G \rightarrow W$ satisfies*

$$\mathcal{E}(f) \leq N^3 - 4N^2 + (10 + 2\kappa)N - (6 + 2\kappa). \quad (\text{F.16.7})$$

Equality holds precisely for an affine map modified at one point by a nonzero value, with that value required to have order two when $\kappa = 1$.

If $G = \mathbb{F}_p^n$ and $W = \mathbb{F}_p^m$, $m, n \geq 1$, the same conclusion holds without $N \geq 34$, provided p is odd or $n \geq 2$. For $p = 2, n = 1$ every map is affine, so there is no non-affine maximum.

[write] **7 October 2026; sources 54–59.** Cutoff lowered for targets without involutions (source 59, Section II.248): the formula with $\kappa = 0$ holds for $N \geq 16$ (its Corollary II.251.3), and its map $\mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/3\mathbb{Z}$ of energy $1296 > 1266$ shows that 16 cannot be lowered below 13; with the table of Remark W2 the exact cutoff for such targets is between 13 and 16 (item II.21). Targets with involutions keep $N \geq 34$.

[write] **7 October 2026; sources 60–69.** The exact cutoff for elementary abelian 3-group targets is 13 (source 67, Section II.293): the formula with $\kappa = 0$ holds for $N \geq 13$ and fails at $N = 12$ (its Corollary II.295.5, with source 59’s $\mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/3\mathbb{Z}$ map). For other targets without involutions the cutoff stays between 13 and 16 (item II.21).

F.16.3 Status, prior literature, and limits

The proofs of the stated theorems are supplied below, including small cases. The correction step is a self-contained version of the classical majority-decoding method associated with Blum, Luby, and Rubinfeld [49]. The derivative-collision interpretation of graph energy and its relation to derivative imbalance and ambiguity are established parts of the literature; see Fu, Feng, Wang, and Carlet [50]. We do not claim these foundational identities or the existence of constant-factor correction as new. The results emphasized here are the explicit small-error cubics, the error-color comparison with remainder terms, the extremizer classification, and the resulting exact gap. A focused literature check does not establish publication priority for these refinements; this article should not be treated as a priority claim.

The results do not improve the global density bound for progression-free sets. For comparison, Leng, Sah, and Sawhney prove for $k \geq 5$ a bound of the form $r_k(N) \ll N \exp(-(\log \log N)^{c_k})$ [2]; our argument does not replace their inverse-theorem or density-increment machinery. No result in this article is advertised as a solution of a previously named open problem in that global theory.

F.17 Source 12. Purpose, main results, and relation to the sources

Source 12, delivered Section 1.

Added in the third write, 6 October 2026. Delivered title block. “Sharp Freiman Extension and Polynomial Bohr Radii / Explicit refinements of Section 10 of Gowers’s proof of Szemerédi’s theorem / Research article and reproducibility supplement / 6 October 2026”; no PDF author field. Archive `gowers_bohr_extension.zip` (arrival 08ab4187e), manuscript `gowers_bohr_extension.tex` (1,215 lines, 25 pages), pin 9ea383e9b. Printed below: the abstract and Section 1; Sections 2–9 and Appendix A form a chapter of Part II (Section II.33), Sections 10–11 and Appendix B are in Part V. Its macro `\snapshot` (its own pin) is renamed `\snapshottwelve`; its `\cleveref` references are written out as `\cleveref` prints them, and its `\tabularx` columns are set ragged right inside its tables, as delivered.

Abstract of source 12 (as delivered). We develop a quantitative replacement for the dense-set extension and Bohr-overlap steps in Section 10 of Gowers’s proof of Szemerédi’s theorem. If $S \subseteq B$ has missing density at most ε and translation by every element of C removes at most a fraction κ of B , then every Freiman homomorphism on S induces a Freiman homomorphism on C provided $3\varepsilon + 2\kappa < 1$. Every rational point of the boundary $3\varepsilon + 2\kappa = 1$ in the relevant parameter range admits an explicit counterexample. We also prove a version with errors measured by collisions of directional increments, higher-order extension statements, and a common-fibre synchronization lemma.

An elementary selection among $4d$ nested Bohr sets replaces the radius $2^{-(d+4)}\rho^d/d$ in Gowers’s Corollary 10.11 by $\rho/(8d)$, after choosing the ambient radius in $[\rho/2, \rho]$. A fourth-moment Fourier estimate then enlarges the initial radius in Theorem 10.13 from $2^{-148}\alpha^{18}/\pi$ to $2^{-74}\alpha^9/\pi$, without changing its large spectrum. For $0 < \alpha \leq 1/6$, selection among $3d$ radii and a careful count of the final slice give a local model on a Bohr set of radius at least $\alpha^{19}/(3\pi 2^{149})$, with retained set of size at least $\alpha^5|X|/20000$. A parameterized form covers all densities $0 < \alpha \leq 1$.

All new estimates are proved here. The application uses explicitly stated, established component-extraction lemmas from Gowers’s argument. Regular Bohr sets, induced difference maps, and local-to-global correction are classical ideas; historical priority for the explicit refinements is not asserted. No improved global bound for Szemerédi’s theorem or new Lean certification is claimed.

The aim is a focused improvement of a complete part of Gowers’s argument, rather than a collection of unrelated constant changes. Section 10 of [1] passes from a multifunction which respects almost all additive relations of order eight to an exact local model on a Bohr neighbourhood determined by the original large spectrum. Two losses in this passage can be removed: a coarse estimate of a Fourier fourth moment, and an unnecessarily small neighbourhood used to guarantee overlap. A third loss is an extra power of the density in the final slicing step.

The point of preserving the original spectrum is structural. Later sections compare spectra arising from many related fibres. Replacing each spectrum by one attached to an independently selected subset would lose those relations. The results below retain this aspect of the argument.

F.17.1 A quantitative overview

For a finite nonempty set B in an abelian group, define its normalized translation loss by

$$b_B(t) = \frac{|B \setminus (B - t)|}{|B|}.$$

The following are the central conclusions; their precise hypotheses appear in the numbered statements.

Sharp extension. If $|S| \geq (1 - \varepsilon)|B|$, $b_B(t) \leq \kappa$ for every $t \in C$, and

$$3\varepsilon + 2\kappa < 1,$$

then a Freiman homomorphism $\phi : S \rightarrow H$ has a well-defined difference map on $C + C$, additive on sums of two elements of C . Its restriction to C is a Freiman homomorphism. There are counterexamples on every rational point of the boundary.

Polynomial local radius. Under the corrected Section 10 hypotheses, with $0 < \alpha \leq 1/6$ and the same spectrum

$$\Gamma = \{\chi : |\widehat{g}(\chi)| \geq 2^{-37}\alpha^{11/2}MN\}, \quad d = |\Gamma|,$$

one obtains an exact local model on

$$C = \mathcal{B}\left(\Gamma, \frac{2^{-74}\alpha^9}{6\pi d}\right), \quad |Y| \geq \frac{\alpha^5}{20000}|X|.$$

Since $d \leq 2^{74}\alpha^{-10}$, its radius is at least

$$\frac{\alpha^{19}}{3\pi 2^{149}}.$$

F.17.2 What is proved, and what is borrowed

The set-theoretic intersection estimates, extension theorems, sharpness examples, error-correction statements, radius-selection lemma, Fourier estimates, and parameter optimization are proved in full. To apply them to Gowers’s multifunction theorem we use the existing component-extraction mechanism from Lemmas 10.3–10.9. Its exact required output is stated in lemma II.40.1; the proof of the application then tracks all subsequent steps and constants.

Quantity	Baseline in the source	Refinement here
Extension from a dense set	Missing density $1/8$ and small prescribed overlap loss	Full sharp range $3\varepsilon + 2\kappa < 1$
Inner Bohr radius	$2^{-(d+4)}\rho^d/d$	$\rho/(8d)$ after ambient-radius selection
Initial spectral radius	$2^{-148}\alpha^{18}/\pi$	$2^{-74}\alpha^9/\pi$ with the same spectrum
Final radius in the sparse-density range	$2^{-155k}\alpha^{18k}/k$	At least $\alpha^{19}/(3\pi 2^{149})$
Retained fraction in corrected repository statement	$\alpha^6/20000$	$\alpha^5/20000$
Density range of the fixed-parameter Lean theorem	$0 < \alpha \leq 1/6$	$0 < \alpha \leq 1$ with a parameterized schedule

Table F.17.1: The radius-selection row changes the ambient-radius quantifier: the dense subset must be selected *after* that radius is chosen. Here $k = \lceil 2^{74}\alpha^{-10} \rceil$ is a rank bound. The retained-fraction comparison is with the corrected repository theorem, not the larger unsupported count in the printed proof.

The relevant repository snapshot is `9ea383e9b28d885a936d5d681cd5584fffd2cb29`. Its `FORMALIZATION_STATUS.txt` reports 97 exact numbered companion proofs and 23 open statements, with 154 modules compiled and 1,508 public theorems audited [66]. *[6 October 2026; sources 11–14] Stale since the pin:* at HEAD `e1db00e99` the ledger records 100 exact companions and 20 open statements (`lemma_16_6`, `lemma_16_9` and `lemma_13_8` have closed since `9ea383e9b`), and `FORMALIZATION_STATUS.txt` reports 206 modules and 1,879 audited public theorems (Section V.1). These are the repository’s recorded verification results. The present work inspected the relevant statements and proof sources, but did not rerun that build or formalize the new theorems.

The repository’s reconstructed L^AT_EX source contains documented corrections to the printed argument [63]. In particular, its safe retained fraction is $\alpha^6/20000$, and the current Lean statement of Theorem 10.13 explicitly assumes $\alpha \leq 1/6$. *[6 October 2026; sources 11–14] Checked at HEAD e1db00e99:* the exported α^6 is a weakening inside the formal proof itself. In `Proofs10Main.lean` the proof of `theorem_10_13_holds` (line 1292) establishes $|W| \geq \varrho^2 \alpha^2 LN/16$ with $\varrho = \alpha^2/32$ (`hregularLower`, line 1398), that is $|W| \geq (\alpha^5/16384)|X|^4$, and $(5/6)|W| \leq |W_1|$ (`hfactor`, `hW1lower`); only then does it spend $\alpha \leq 1/6$ (`halphaScaled`, line 1407) to export $\alpha^6/20000$. The docstring (`Section10.lean`, lines 467–472) and the transcription’s editorial note call $\alpha^6/20000$ the safe, proof-supported conclusion; that is conservative, not wrong. We do not treat a source heading or an informal displayed formula as a substitute for checking these hypotheses.

F.17.3 Literature and novelty

The advantage of regular Bohr sets over Gowers’s original overlap estimate is established in the literature. Green and Tao’s inverse theorem [4], Definition 2.6 and Lemma 8.2, gives the standard regular-radius framework; its discussion explicitly relates this to Gowers’s Lemma 10.10. We use a simpler finite selection argument because only one overlap scale is required. The induced difference map is also standard, and concentrated increment distributions appear in work on Freiman extension, including Conlon and Gowers [51].

Thus the contribution of this article is an explicit sharp extension criterion and its complete quantitative integration into this particular Section 10 parameter chain, together with reusable generalizations. It is not a claim to have invented Bohr regularization or to have superseded modern

inverse theorems. Modern quantitative Szemerédi results already go beyond the original global bound: for example, Leng, Sah, and Sawhney prove $r_\ell(N) \ll N \exp[-(\log \log N)^{c_\ell}]$ for $\ell \geq 5$ [2]. The conclusions here concern a local structural step.

F.18 Source 13. Scope, results, and the source baseline

Source 13, delivered Section 1.

Added in the third write, 6 October 2026. Delivered title block. “Research article / Sharp Local Refinements in Gowers’ Szemerédi Framework / Simultaneous recurrence, extremal density selection, and centered cube counts / Prepared for Vladimir Reshetnikov / 6 October 2026”; the PDF author field reads “Research report prepared for Vladimir Reshetnikov”. Archive `gowers_local_refinements.zip` of arrival `08ab4187e`: *a new archive that reuses the name of source 03’s archive* (and its manuscript reuses source 03’s file name); it is an independent manuscript, not an edition of source 03 (different title, pin and files; 0.97% shared text), and it is retrieved with `git show 08ab4187e:docs/incoming/gowers_local_refinements.zip` (source 03’s with `66f24b0d0`). Manuscript `gowers_local_refinements.tex` (1,223 lines, 31 pages), pin `9ea383e9b` (its `provenance.json`). Printed below: the abstract, the title page’s snapshot line and Section 1; Sections 2–6 are in Part I, 7–9 in Part III, 10–11 and Appendix A in Part V. Its ‘question’ environment, printed “Research question” there, is set in this report’s ‘Research question’ style.

Abstract of source 13 (as delivered). We develop explicit refinements of several local ingredients in Gowers’ proof of Szemerédi’s theorem, using the supplied ProveIt paper and Lean development as the comparison target. For q linear phases, a simultaneous approximation argument replaces the exponent $1/(2 \cdot 16^q)$ in the linear specialization of the prescribed-length partition lemma by $1/(3q + 2)$, improves the diameter bound simultaneously, and removes its large-size threshold. For flexible cell lengths we prove the optimal worst-case exponent frontier $(q + 1)a + qb = 1$, together with finite modular obstruction certificates. We then solve two extremal selection problems: conversion of partition imbalance into density on a large cell, and joint conversion of phase correlation, phase diameter, and exceptional mass into a high-density-cell mass bound. These results give a density-sensitive Fourier increment with an explicit progression length. A separate refinement exploits zero-mean cancellation in cube expansions: the linear, quadratic, and cubic perturbation terms vanish, the quartic term is counted exactly, and the three-dimensional cube admits a complete spectral decomposition. Finally, we retain interval length in Fourier discrepancy and telescope progression counts to improve a 2^k error factor to a density-weighted sum of at most $k - 2$ terms. All principal claims have proofs; executable checks and a formalization roadmap are supplied. The report establishes local mathematical refinements and sharpness statements, without asserting historical priority or a new global record for Szemerédi’s theorem.

Source snapshot: ProveIt, commit `9ea383e9b28d885a936d5d681cd5584fffd2cb29`.

Gowers’ proof organizes a long argument around a small number of recurring operations: recognize uniformity, extract algebraic structure from failure of uniformity, partition that structure into pieces on which phases vary little, turn correlation into a density increment, and iterate. Several individual bounds deliberately sacrifice considerable information. This article improves selected local operations while keeping their hypotheses and outputs explicit.

The comparison is with two distinct sources. The journal article is Gowers’ *A new proof of Szemerédi’s theorem*, published in 2001 [1]. The user-supplied ProveIt tree contains an adapted L^AT_EX transcription and a Lean development [67, 68]. The adapted text includes editorial repairs to hypotheses, numerical constants, and rounding. Throughout, “the repository version” refers to

commit 9ea383e9b28d885a936d5d681cd5584fffd2cb29, whose commit timestamp is 2026-10-06 20:35:24 UTC. Statements below are compared to that fixed snapshot, rather than to a moving branch.

F.18.1 What is proved

The most substantial local conclusions are as follows.

- A. Simultaneous linear partitions.** For q affine maps to $\mathbb{Z}/N\mathbb{Z}$, put $a_q = 1/(3q + 2)$. For every $1 \leq r \leq N$ and every integer $1 \leq v \leq r^{a_q}$, there is a complete partition of $\{0, \dots, r - 1\}$ into nonempty arithmetic progressions of lengths $v - 1$ or v , on each of which every image has modular diameter at most $r^{-2a_q}N$. This directly strengthens the linear specialization of the repository’s Lemma 5.9. The proof is finite and has no asymptotic threshold.
- B. An optimal exponent tradeoff.** When lengths may range from L to $2L - 1$, a partition with $L \asymp r^a$ and simultaneous phase diameter bounds of order r^{-b} is universally possible at the frontier $(q + 1)a + qb = 1$. A random-frequency obstruction proves that no larger exponent region is possible. The sharpness assertion concerns this flexible-length problem.
- C. Sharp density and phase selection.** A bounded mean-zero function carries more useful information than its absolute bound alone. We solve the weighted extremal problem for high-average cells, including discarded mass, and then solve the joint problem with a phase diameter constraint. Explicit extremizers show that the mass bounds cannot be improved using only the stated data.
- D. Centered cube counts.** Over $\mathbb{F}_p^n$ with p odd, balancedness annihilates many terms in a cube expansion. The quartic coefficient is

$$P_d = \frac{6^d - 2 \cdot 4^d + 2^d}{8},$$

and multiplies $\delta^{2^d-4} \|f\|_{U^2}^4$. For $d = 3$, the surviving subset counts are

$$1, 0, 0, 0, 12, 8, 28, 8, 1.$$

We give a complete identity and stronger bounds in terms of the Fourier spectrum.

- E. Two additional quantitative refinements.** Interval discrepancy gains the sharp length exponent $\min(\beta, 1 - \beta)^{1/4}$, and telescoping replaces the 2^k error factor in the progression-counting corollary by $\sum_{j=3}^k \prod_{i>j} \delta_i \leq k - 2$ under the same uniformity budget.

F.18.2 Novelty and mathematical status

The simultaneous approximation, Fourier duality, Gowers–Cauchy–Schwarz, and convexity tools used here are classical. The report’s contributions are the precise strengthened statements, exact finite rounding, extremal tradeoffs, spectral decomposition, and their integration with the named lemmas. A search of primary sources did not establish historical priority for the exact package of results. Accordingly, “proved here” is a claim about the arguments in this article, not a claim that no earlier author has derived a related inequality.

Nor do these local improvements imply a new best bound for the largest progression-free subset of $[N]$. For comparison, Leng, Sah, and Sawhney proved in 2024 that for every $k \geq 5$,

$$r_k(N) \ll N \exp(-(\log \log N)^{c_k})$$

for a positive constant c_k [2]. Improving a linear subroutine in the 2001 proof does not bypass its higher-degree structure theorem. We explicitly identify the additional work needed to obtain a global improvement in section V.54.

F.18.3 A map to the supplied development

Source target	Information retained here	Resulting improvement
Lemma 5.9, $k = 1$	Simultaneous frequency dimension; integer cell decomposition	Exponent $1/(3q + 2)$ and smaller diameter; no giant threshold
Lemmas 5.14–5.15	Signs, ambient density, phase chord diameter, exceptional mass	Sharp weighted selection and a sharp joint frontier
Corollary 2.5 Lemma 3.10	Exact $\mathbb{E} \mathbf{1}_A - \delta = 2\delta(1 - \delta)$ Mean-zero cancellation and affine dependencies of cube vertices	Length gain proportional to $[\delta(1 - \delta)]^{-1/2}$ Exact quartic coefficient and a full 3-cube decomposition
Lemma 3.5	Both Fourier bounds: interval length and frequency decay	Sharp quarter-power dependence on interval length
Corollary 3.3	Telescoping and exact two-form independence	Density-weighted error instead of 2^k terms

The repository source ledger reports 97 numbered statements with exact companion declarations and 23 open entries. *[6 October 2026; sources 11–14] Stale since the pin* (also in this source’s `provenance.json`): at HEAD `e1db00e99` the ledger records 100 and 20 (`lemma_16_6`, `lemma_16_9` and `lemma_13_8` have closed since `9ea383e9b`; Section V.1). Its own scope warning says that this is a source inventory, not certification of compilation or fidelity to the journal paper [69]. We use the Lean files to identify interfaces and normalization. No new Lean compilation or kernel verification is claimed in this report; the new proofs are written mathematical proofs, supplemented by finite computations.

F.19 Source 14. Scope, contributions, and relation to the sources

Source 14, delivered Section 1.

Added in the third write, 6 October 2026. Delivered title block. “Additive combinatorics / Sharp Extremal Refinements in Gowers’s Szemerédi Framework / Secondary cube asymptotics, arrangement stability, and polynomial supersaturation / Prepared for Vladimir Reshetnikov / Research report prepared by ChatGPT / 6 October 2026”; *the PDF author field reads “ChatGPT, prepared for Vladimir Reshetnikov”*. Archive `gowers_extremal_refinements.zip` (arrival `d179062cc`), manuscript `gowers_extremal_refinements.tex` (2,767 lines, 37 pages), pin `9ea383e9b` (its `source_manifest.json`). Printed below: the abstract, the title page’s statement of principal conclusions and its closing note, and Section 1; Sections 2, 3, 6, 7 and Appendix A are in Part III, 4–5 in Part II, 8–9 and Appendix B in Part V. Its macro `\leanfile` is kept; its `\repo` (unused) is dropped.

Abstract of source 14 (as delivered). We study three local extremal problems suggested by Gowers’s proof of Szemerédi’s theorem and its accompanying formalization. First, for a real mean-zero function f on a finite abelian group of order coprime to six, with $u = \|f\|_{U^3}$, we prove the uniform upper expansion

$$\|\delta + f\|_{U^3}^8 \leq \delta^8 + 6\sqrt{2}\delta^4 u^4 + (9/4)^{1/3}\delta^{8/3}u^{16/3} + O(\delta^{7/3}u^{17/3}).$$

The secondary coefficient is attained asymptotically by a precisely tuned two-harmonic weight, and near attainment forces the same amplitude, phase alignment, and Fourier concentration. Additional order-five resonances change the next extremal exponent from $16/3$ to 6 on groups of exponent five.

Second, if a set of size m supports at least $(1 - \varepsilon)m^3$ respected pairs of congruent coordinate cubes, with $0 \leq \varepsilon \leq 1/100$, it is within $\frac{1}{2}(1 - \sqrt{1 - 4\varepsilon})m$ points of a unique product of subgroup cosets. An infinite family attains this error envelope. We classify equality in the single-section collision bound, the arrangement moment inequality, and their combination. Geometric defects give an explicit stronger denominator in higher arrangement counts.

Third, for $k = 2m \geq 4$, characteristic $p > k$, and polynomial color coordinates of degrees at most $D_\ell \leq k$, put $B = \sum_\ell [D_\ell/2]^2$. The exact minimum number of ordered nontrivial symmetrically colored k -term progressions on $\mathbb{F}_q^d$ is $q^d(q^{\max(d-B,0)} - 1)$. Equality forces a linear space of directions at every center, and the first supercritical prime-field dimension has a sharp gap. Classical minimum-support theorems and norm constructions are explicitly credited. The conclusions refine local components and solve specified finite extremal questions; they do not supply a new global quantitative Szemerédi bound.

Three principal conclusions. The optimal secondary coefficient in the centered three-cube problem is $(9/4)^{1/3}$; near-maximal congruent-cube collisions force a unique nearby product of subgroup cosets with error $(1 - \sqrt{1 - 4\varepsilon})/2$; and the exact number of unavoidable symmetric arithmetic progressions is determined for polynomial colorings with specified low coordinate degrees.

Based on the supplied Gowers paper and the accompanying Lean development, inspected at repository commit `9ea383e9b28d885a936d5d681cd5584fffd2cb29`. The report includes complete written proofs, exact finite certificates, reproducible diagnostics, a source-to-theorem map, and a research agenda. The results have not been externally peer reviewed or newly certified by the Lean kernel; bibliographic priority is not asserted.

Gowers’s proof [1] makes repeated use of a productive principle: a configuration count is controlled by a norm or moment, and a sufficiently large value eventually yields algebraic structure. The coarse estimate is often enough to finish the density-increment argument. This leaves several worthwhile local questions. What is the correct first or second term of an extremal count? Which examples attain a moment inequality? How much structure survives near equality? Can an existence obstruction be upgraded to an exact counting theorem?

This report gives three connected answers. Its scope is the precisely stated extremal models below. Every main claim comes with a written proof, including the finite-field support lemmas used in the coloring application. The supporting computations check delicate finite identities and small equality cases; the general claims rest on the proofs.

F.19.1 What is strengthened

1. **Three-cube counts.** Gowers’s Lemma 3.10 bounds the normalized number of d -cubes by $(\delta + u)^{2^d}$, where u is the normalized uniformity norm of the balanced function. Centering

removes the apparent first-order error when $d = 3$. Theorem III.34.1 goes further: it determines the optimal coefficient of the next fractional power after the sharp quartic term. Theorem III.34.3 identifies the associated extremizing Fourier geometry.

2. **Collision and arrangement counts.** The moment step behind Lemmas 12.3 and 14.7 can retain the cross-section profile $S = \sum_r |B_r|^3$. Theorems II.43.2 and II.43.3 identify what almost saturating that profile means. Theorem II.44.4 classifies simultaneous equality in the geometric and spectral steps. Equation (II.44.4) improves the moment denominator whenever a geometric defect is known.
3. **A related finite-field coloring question.** Theorem III.36.1 determines an exact minimum in a restricted but natural polynomial-coordinate model. It strengthens an existence threshold to pointwise and global supersaturation, with equality geometry and an arithmetic gap. The symmetry condition is defined explicitly; it differs from the monochromatic condition in Gowers’s Corollary 8.3.

An earlier project note [58] established the sharp quartic cube coefficient $6\sqrt{2}$, exhibited the fractional $u^{16/3}$ correction, and obtained an upper secondary coefficient 6 against a lower example $(9/4)^{1/3}$. A separate project report [59] established the polynomial-coloring dimension threshold. The present report closes the cube coefficient gap and determines the full minimum count above the coloring threshold. All prerequisites from those unpublished reports that are needed here are reproved, so the proofs can be read independently.

F.19.2 Classical ingredients and limits of the comparison

The real, centered comparison $\|f\|_{U^3}^8 \geq 2\|f\|_{U^2}^8$ leads to a single Fourier pair at leading order. Our secondary problem measures a small violation of that equality structure. It is different from comparing a Gowers norm to an L^p or L^∞ norm and classifying values near their maximum; the latter is treated in depth by Eisner and Tao [7]. We use elementary finite Fourier analysis for the particular comparison needed here.

Near-maximal additive energy and proximity to a coset are classical. For example, Proposition 1.1 in Sanders’s survey [54] displays a conclusion with square-root loss in the normalized energy defect. The elementary disjoint-defect count below yields the explicit envelope $\tau(\varepsilon) = \varepsilon + \varepsilon^2 + 2\varepsilon^3 + \dots$ in a specified small-defect range, and then forces coordinate splitting for cube collisions. This is a comparison with that displayed proposition, not a claim that no linear stability estimate has appeared elsewhere. Karam [55] discusses related coset-approximation questions under different distributional small-doubling hypotheses; those assumptions should not be identified with ours.

Warning’s second theorem [40, 56] supplies a lower bound for a nonempty finite-field zero set from the sum of degrees. The equality case used here follows from the classical minimum-weight Reed–Muller theorem, for which Leducq [57] gives a modern proof. Odd norm layers already occur in the symmetric-coloring construction of Shi and Dong [33]. Our counting and rigidity conclusions combine these ingredients after an exact reduction to odd Taylor coefficients. The ingredients themselves are not new results of this report.

The local results must also be distinguished from an improvement of the final density bound for long progressions. For context, Leng, Sah, and Sawhney [2] obtained, for every $k \geq 5$, a bound of the form

$$r_k(N) \ll_k N \exp(-(\log \log N)^{c_k}), \quad c_k > 0,$$

where $r_k(N)$ is the largest size of a subset of $\{1, \dots, N\}$ with no nontrivial k -term arithmetic progression. No new bound for $r_k(N)$ is derived here. The research questions at the end identify which additional bridges would be needed to turn the local gains into a global improvement.

F.19.3 Conventions that matter

There are two different kinds of cube in this report. A three-dimensional *additive cube* in one abelian group has parameters $(x, h_1, h_2, h_3) \in G^4$ and eight vertices $x + \omega_1 h_1 + \omega_2 h_2 + \omega_3 h_3$. Its normalized weighted count is $Q_3(W)$. An *axis-parallel coordinate cube* in $G_1 \times \dots \times G_k$ has a base x and a vector of coordinate sidelengths h ; its vertices choose either x_i or $x_i + h_i$ in each coordinate. Congruent pairs of these cubes are counted by $Q_\phi(A)$. The distinction prevents a mistaken identification of their parameter spaces or normalizations.

All averages $\mathbb{E}$ are uniform probability averages. Fourier transforms in the cube section are normalized by $|G|^{-1}$. The arrangement section explicitly uses an unnormalized transform, with the compensating $|R \times K|^{-1}$ in its moment formula. All combinatorial counts are counts of ordered parameter tuples. Zero sidelengths and repeated cube vertices are included. Nontrivial arithmetic progressions in the coloring section have nonzero direction, and their points are distinct under $p > k$.

The constant in $O_\delta(\cdot)$ may depend on a fixed density δ . An unadorned $O(\cdot)$ in the normalized cube proof has an absolute constant, independent of both the group and the function. The restrictions on order or characteristic are part of each statement and are not silently omitted in applications.

F.19.4 An explicit translation back to Lemma 3.10

Let $A \subset \mathbb{Z}_N$, let $\delta = |A|/N > 0$, assume $\gcd(N, 6) = 1$, and write $f = 1_A - \delta$. In the normalization of the paper and the inspected Lean statement, degree-two α -uniformity means $\|f\|_{U^3}^8 \leq \alpha$. Theorem III.34.1 therefore gives, whenever $\alpha^{1/8} \leq c\delta$,

$$\begin{aligned} \#\{\text{three-cubes in } A\} &\leq N^4 \left(\delta^8 + 6\sqrt{2} \delta^4 \alpha^{1/2} + (9/4)^{1/3} \delta^{8/3} \alpha^{2/3} \right. \\ &\quad \left. + C \delta^{7/3} \alpha^{17/24} \right). \end{aligned} \tag{F.19.1}$$

The passage from the actual norm to its budget uses monotonicity of the positive powers in the upper bound. At fixed positive density, the error begins at $\alpha^{1/2}$, whereas expanding the bound in Lemma 3.10 begins at $\alpha^{1/8}$. The sharper statement uses the additional mean-zero, real-valued, and coprime-order structure specified above. Its optimal secondary constant is meaningful asymptotically; it is not a claim that the displayed truncated expression without its remainder is a finite-parameter bound.

F.20 Source 15. Results, scope, and the source discrepancy

Source 15, delivered Section 1.

Added in the fourth write, 6 October 2026. Delivered title block. “Collision-Sensitive Amplification of Respected Arrangements / Sharp refinements of Gowers’s Lemma 14.7, coupled product-property bounds, and Cartesian stability / Research manuscript prepared for Vladimir Reshetnikov / 6 October 2026”; the PDF author field reads “Research manuscript prepared for Vladimir Reshetnikov” and the running head “Research manuscript”. Archive `collision_sensitive_arrangements.zip` (arrival 212999d4b), manuscript `collision_sensitive_arrangements.tex` (1,752 lines, 27 pages), pin 9ea383e9b (its `provenance.json`, shipped as `data/15-collision-amplification-provenance.json`). Printed below: the abstract and Section 1; Sections 2–12 form a chapter of Part II (Section II.45), Sections 13–15 and Appendices A–B are in Part V.

Abstract of source 15 (as delivered). We retain the exact collision count that is discarded in the second-moment estimate of Gowers’s arrangement-amplification argument. If R_d counts ordered respected d -arrangements and $C = R_1$, then $R_d C^{d-2} \geq R_2^{d-1}$ for every integer $d \geq 2$. The collision count satisfies the sharper geometric bound $C \leq \sum_x |B_x|^3$, giving sharp density-sensitive versions of Lemma 14.7. More substantially, retaining the same parameter in the preceding product-property step gives $a_2 \geq \gamma^{8 \cdot 2^k} \chi^4$ and hence $a_d \geq \gamma^{8(d-1)2^k} \chi^{3d-2}$, in normalized notation. For eight-arrangements this yields $\beta^{22 \cdot 4^k} \gamma^{44k4^k + 56 \cdot 2^k}$, strengthening the repository’s corrected product-property bound. We also characterize equality in the cubic collision bound and prove an explicit dimension-free stability result: a set with at least $(1 - \varepsilon)|A|^3$ parallel axis-cube pairs is within $18\varepsilon|A|$ points of a Cartesian coset when $\varepsilon \leq 1/100$. Further results include spectral stability, side-vector refinements, weighted extensions, sharp examples, and an integration plan for the supplied Lean development. All new assertions are proved in ordinary mathematics; finite tests are supplied separately. No improvement to a final Szemerédi bound, independently checked Lean formalization, or literature-priority claim is made.

Gowers’s Section 14 converts a supply of respected pairs of parallelepipeds into a supply of longer respected arrangements. The relevant analytic step is a moment interpolation, while its quantitative input comes from a weighted simultaneous-additivity property. These two steps share a natural intermediate parameter. Eliminating it independently in each step loses information. This manuscript instead keeps that parameter until the steps have been combined.

The source comparison uses the paper [1] and the supplied ProveIt files at commit 9ea383e9b28d [78, 79, 80, 82]. The finite cyclic notation in the paper is generalized here to arbitrary finite abelian coordinate groups and an arbitrary finite abelian value group. Every count is ordered and allows repeated vertices. These conventions are essential to the exact identities.

F.20.1 Three levels of improvement

Write R_d for the respected d -arrangement count, $C = R_1$ for the collision count, and B_x for a horizontal fibre. The first improvement is

$$R_8 \geq \frac{R_2^7}{C^6} \geq \frac{R_2^7}{(\sum_x |B_x|^3)^6}. \quad (\text{F.20.1})$$

The first inequality is the exact form of the interpolation already underlying the paper; the second retains the geometry of the domain instead of replacing it by the entire ambient space. The constant is sharp.

The second improvement couples this interpolation to the preceding product-property argument. For horizontal ambient size M , vertical size n , and density $\beta = |B|/(Mn)$, set

$$a_d = \frac{R_d}{M^{2d+1}n^{2d-1}}, \quad \chi = \frac{C}{M^3n}.$$

Under the product property with parameter γ , we prove

$$a_d \geq \gamma^{8(d-1)2^k} \chi^{3d-2} \geq \beta^{(3d-2)4^k} \gamma^{(6d-4)k4^k+8(d-1)2^k}. \quad (\text{F.20.2})$$

In particular,

$$\boxed{a_8 \geq \beta^{22 \cdot 4^k} \gamma^{44k4^k+56 \cdot 2^k}}. \quad (\text{F.20.3})$$

Section II.49 supplies the complete proof, including the configuration-count input. No inverse theorem or unproved catalogue statement is assumed in that proof.

The third improvement is structural. For $A \subseteq G_1 \times \cdots \times G_k$, let $Q(A)$ be the number of pairs of contained axis cubes with the same side vector. We prove

$$Q(A) \geq (1 - \varepsilon)|A|^3, \quad 0 \leq \varepsilon \leq \frac{1}{100} \implies |A \triangle (c + H_1 \times \cdots \times H_k)| \leq 18\varepsilon|A|. \quad (\text{F.20.4})$$

The constant does not depend on the number or sizes of the coordinate groups. The linear order in ε cannot be improved uniformly to $o(\varepsilon)$. This answers the specific near-extremal Cartesian inverse question formulated by the left-hand side of (F.20.4). It is a near-maximal-energy theorem, not a sparse inverse theorem.

F.20.2 The printed and corrected product-property exponents

A source-fidelity issue materially changes the comparison. The printed Corollary 14.6 uses $k4^{k+1}$ as the exponent of γ . The supplied definition file explicitly records that its proof only gives

$$E_k = 2k4^{k+1} + 8 \cdot 2^k, \quad (\text{F.20.5})$$

and adopts the safe weakening $3k4^{k+1}$ for $k \geq 1$ and $0 < \gamma \leq 1$. Its Lemma 14.8 therefore has exponent $21k4^{k+1}$, rather than the printed $7k4^{k+1}$ [79, 82]. We reproduce the unweakened exponent (F.20.5) below and use the corrected statement as the benchmark. The discrepancy was already documented in the repository; it is not a discovery claimed by this manuscript.

[write] **6 October 2026; sources 15–20.** *Formal status understated (checked at HEAD 0b4890a1c, unchanged since the pin).* The unweakened exponent (F.20.5) is not only documented: it is the conclusion of the private, kernel-checked lemma `cor146_direct_lower` (`Proofs14Arrangements.lean`, line 759), which proves $\beta^{4^{k+1}} \gamma^{2k4^{k+1}+8 \cdot 2^k} N^{5k+3} \leq R_2$ from the fourth-moment bound of `cor146_average_fourth_moment`. The public `corollary_14_6_holds` (line 838) obtains it and only then weakens it to $3k4^{k+1}$ through the private `cor146_gamma_exponent_1e` (line 817). So the corrected exponent $2k4^{k+1} + 8 \cdot 2^k$ is a proof-level fact of the formal development; the public statement `corollary_14_6` (`Sections14_15.lean`, line 266) has the weakened $3k4^{k+1}$, and `lemma_14_8` (line 292) the resulting $21k4^{k+1}$.

Consequently, (F.20.3) strengthens the *corrected, supported* product-property bound. It must not be described as proving the smaller unsupported printed γ exponent. Lemma 14.7 itself has no such exponent issue, and its refinement (F.20.1) is unconditional.

F.20.3 What is, and is not, claimed

The proofs below establish the displayed finite theorems. The Fourier identities, interpolation mechanism, and configuration induction are established ingredients, not new methods. The contribution developed here is the parameter-retaining combination, its sharp density refinements, and

the explicit quantitative Cartesian analysis. A targeted literature check does not establish priority over all existing work on additive energy, near-extremizers, or product-space inequalities.

[write] **6 October 2026; sources 15–20.** *Credit and novelty within this report.* The pin 9ea383e9b (13:35 on 6 October 2026) predates the placement of sources 01–14 in this directory, so this source could see neither source 04 nor source 14; its results are independent. In this report, however, two of the three contributions named above are not new. The density-sensitive interpolation (F.20.1) with the cubic cap and the balanced factor is source 04’s Theorem II.10.1, displays (A5)–(A11). The “explicit quantitative Cartesian analysis” is superseded by source 14’s Section II.43: under the same hypothesis $Q(A) \geq (1 - \varepsilon)|A|^3$, $\varepsilon \leq 1/100$, source 14’s Theorem II.43.3 gives a unique closest product coset at distance at most $\tau(\varepsilon)|A|$, $\tau(\varepsilon) = (1 - \sqrt{1 - 4\varepsilon})/2 \leq 1.0102\varepsilon$, in place of $18\varepsilon|A|$ (note after Theorem II.52.2). New in this report: the collision-parametrized product step and the coupled bounds (Theorems II.49.5 and II.49.6), the automatic weak parameter (Lemma II.49.2), the band and profile results (Corollary II.53.2, Proposition II.53.3, Example II.53.4) and the vertex weights (Theorem II.56.1). The sentence “This answers the specific near-extremal Cartesian inverse question” above refers to the source’s own formulation; source 14 answers it with the sharp constant.

These are local improvements within an older proof architecture. For comparison, Leng, Sah, and Sawhney prove for $\ell \geq 5$ the global estimate $r_\ell(N) \ll N \exp(-(\log \log N)^{c_\ell})$ [2]. We do not propagate the present changes through every later extraction, recurrence, and density-increment step, and claim no improvement to that or any other final Szemerédi estimate.

F.21 Source 16. Scope, source comparison, and principal results

Source 16, delivered Section 1.

Added in the fourth write, 6 October 2026. Delivered title block. “Local Phase Removal Without Carry Patterns / and further quantitative refinements of Gowers’ proof of Szemerédi’s theorem / Research report prepared for Vladimir Reshetnikov / 6 October 2026”; the same in the PDF author field. Archive `Gowers_Localization_and_Quantitative_Refinements.zip` (arrival 99f28708e), manuscript `Gowers_Localization_and_Quantitative_Refinements.tex` (2,127 lines, 30 pages), pin dca9f0638 (its delivery README and Section 1). Printed below: the abstract and Section 1; Sections 2–5 form a chapter of Part IV (Section IV.10), Sections 6–7 are in Part I, Section 8 in Part III, Sections 9–12 and Appendix A in Part V. Its macros `\conj` and `\BB` are renamed `\conjsixteen` and `\BBsixteen` (the report’s own have other meanings); its equation numbers are delivered as explicit tags $(k.j)$ and are kept, so they print as delivered.

Abstract of source 16 (as delivered). We give three families of local refinements of the arguments in Gowers’ 2001 proof of Szemerédi’s theorem. The principal result replaces the carry-pattern decomposition in Proposition 17.7 by a finite family of sliding interval windows. The resulting relative uniformity factor is

$$\alpha \left(\frac{4}{3(k+2)} \right)^{k+2},$$

in place of the displayed factor $2^{-2(k+1)^3} \alpha$. We retain an exact partition and the required minimum cell length, and give an explicit upper cell length $3km$. A more precise finite formula admits a width optimization whose fixed-dimension limiting coefficient is $(2/(k+2))^{k+2}$. A local-scale obstruction proves that the leading logarithmic loss $k \log k$ is optimal. We also give a version that preserves the original smaller cell scale, replacing a squared carry-count loss by an active-pattern count times a global count.

Second, we construct exact partitions on which affine frequencies are small after multiplication by each cell’s own step. An elementary certificate theorem, combined with a precisely cited simultaneous quadratic recurrence theorem, yields every balanced exponent $\gamma < ((q+1)(42q+1))^{-1}$. A matching problem is formulated through a rigorous obstruction $\gamma \leq (3q+1)^{-1}$, leaving

a substantial and explicit gap. Third, we compute the fourth-order coefficient of centered cubes in every dimension and the fifth-order coefficient over odd-order groups, including the first additional 3-torsion term in dimension four. The resulting error exponent $4/2^d$ is proved optimal even for genuine sets of exact density $1/3$ in odd cyclic groups.

All new assertions are supplied with proofs. Exact combinatorial and rational computations, and a separate numerical localization check, are included for reproducibility. The modern recurrence estimate is an external input. Historical priority of the formulations developed here is not asserted, and the results do not establish a new global Szemerédi bound or a Lean formalization of these refinements.

The target is a collection of concrete, verifiable improvements to local arguments in [1], rather than a claim that the full quantitative Szemerédi problem has been settled. The distinction matters: a substantial improvement at one stage may be numerically dominated by a different stage, or may change a scale that must be transported through the induction.

The source repository was inspected at commit

`dca9f06383ccdfc2ed208ccb7c9882214fdb09ac`.

We used its paper transcription, its Section 17 and Section 13 Lean interfaces, its formalization status, and its existing local-refinements article [84]. The existing research already covers numerous three-dimensional cube estimates, density-selection inequalities, Freiman restrictions, and carry-pattern counts. The present work concentrates on a different localization argument, direct control of moving frequencies on their own steps, and coefficients for cubes of arbitrary dimension.

[write] 6 October 2026; sources 15–20. *What the “existing local-refinements article” was at the pin.* At `dca9f0638` the file `local-quantitative-refinements/article.tex` was source 04’s delivered manuscript (staged by `aefb0b44a`; the first write of this report, `18507e2b2`, came later). It contains source 04’s Section 8 (here Section IV.6), whose Lemmas IV.6.5–IV.6.6 and the first inequality of Theorem IV.6.7 are this source’s Lemma IV.14.1 and Proposition IV.14.2, and its B_5 bound and set sharpness (half of this source’s Proposition III.39.8; Proposition III.39.10). This source does not credit them (notes there). Its Section III.39 is an independent third derivation of source 09’s Theorem III.21.1 (source 09’s manuscript was never staged, so this source could not see it); the “first additional 3-torsion term in dimension four” of the abstract is source 09’s $S_4 = 16$.

F.21.1 A strengthened local phase-removal theorem

Let N be prime, let $f : \mathbb{Z}/N\mathbb{Z} \rightarrow \mathbb{C}$ be bounded in modulus by one, and let

$$P = \prod_{i=1}^k (r_i + d[-h, h]), \quad m = 2h + 1, \quad m^2 \leq N,$$

where $d \neq 0$. For a k -tuple a , write

$$\Delta(f; a)(s) = \prod_{\epsilon \in \{0,1\}^k} C^{|\epsilon|} f(s - \epsilon \cdot a), \quad \hat{u}(r) = \sum_s u(s) e_N(-rs), \quad e_N(t) = e^{2\pi i t/N}.$$

A polynomial is called *multi-affine* if its degree in each variable is at most one. This includes the affine lower-degree terms that occur after translating a multilinear expression.

Theorem F.21.1 (Main phase-removal consequence). *Suppose $k \geq 1$, $\alpha > 0$, and σ is the restriction to P of a multi-affine polynomial over $\mathbb{F}_N$. If*

$$\sum_{a \in P} \left| \widehat{\Delta(f; a)}(\sigma(a)) \right|^2 \geq \alpha N^2 m^k, \quad (1.1)$$

then there are a polynomial $\phi \in \mathbb{F}_N[X]$ of degree at most $k + 1$, an integer $l \geq 1$, and a partition $\mathcal{Q}$ of $\mathbb{Z}/N\mathbb{Z}$ into proper modular arithmetic progressions, all of cardinality l or $l + 1$, such that

$$l \geq \frac{m}{3k} \quad \text{and} \quad \frac{\sum_{Q \in \mathcal{Q}} \mathcal{C}_{k+1}(\mathbf{1}_Q f e_N(-\phi))}{|\mathcal{Q}|(l+1)^{k+2}} > \alpha \left(\frac{4}{3(k+2)} \right)^{k+2}. \quad (1.2)$$

Here $\mathcal{C}_j(u) = N^{j+1} \|u\|_{U^j}^{2j}$ is the unnormalized j -cube form. When $m > 3k$, the construction gives the stronger lower bound $l \geq m/3$, and $l + 1 \leq 3km$.

The proof occupies Sections IV.12–IV.13. It changes the geometry of the intermediate cells. Gowers' proof uses relatively short cells and counts all the ways a cube can cross their boundaries; we average windows that contain the entire cube and count their coverage multiplicity exactly. The upper size $3km$ is therefore part of the result and part of any downstream accounting.

The quantity in (1.2) is precisely the relevant partition-uniformity quotient with upper size $l + 1$. Thus the strict inequality states failure of uniformity at the displayed factor. The comparison with Proposition 17.7 of [1] is at this explicit normalization, rather than at an unstated choice of cell size.

[write] **6 October 2026; sources 15–20.** *Theorem F.21.1 strengthens a statement that is already formal* (checked at HEAD 0b4890a1c; the source does not say so). Proposition 17.7 is `proposition_17_7` (Sections17_18.lean, line 140) with the exact companion `proposition_17_7_holds` (Proofs17LocalizedPhaseRemoval.lean, line 2900, proved since e336d1b3a, 24 August 2026, and present at the pin). Its hypotheses are those of the theorem here: N prime, $0 < \alpha$, $0 < k$, m odd with $m \leq \sqrt{N}$, a box of width m with nonzero common difference, a multilinear σ , and $\alpha N^2 m^k$ at most the sum of the squared Fourier coefficients (with `cubeDifference` and `fourier` the conventions here). Its conclusion is a polynomial phase of degree at most $k + 1$ and a partition into proper modular progressions of lengths l or $l + 1$ with $m/(3k) \leq l$ on which `UniformOnPartition` (line 132), normalized by $(l + 1)^{k+2}$, fails at $2^{-2(k+1)^3} \alpha$; that is exactly (1.2) with a smaller constant. Since $c_k = (4/(3(k+2)))^{k+2} > 2^{-2(k+1)^3}$ for every $k \geq 1$ (proof at the write: $\log_2(3(k+2)/4) < k+2$, so $\log_2(1/c_k) < (k+2)^2 \leq 2(k+1)^3$), the theorem implies the formal statement, with a far better constant (29.07 against 1,458 lost bits at $k = 8$; Section IV.10). It is a drop-in strengthening of a proved statement; it is not formalized. The output bound $l + 1 \leq 3km$ is information the formal statement does not ask for. Under the printed reading of the cell size as “at most m ” (source 04’s), these cells are too long; the formal statement and source 06 use the $(l + 1)$ normalization, as Section 18 does. No formal declaration consumes `proposition_17_7` (`theorem_18_1` is open).

F.21.2 Direct control of the step

For affine frequencies $A_i n + B_i \in \mathbb{R}/\mathbb{Z}$, we construct a partition of $\{0, \dots, R - 1\}$ into progressions Q satisfying

$$\|(A_i n + B_i) d_Q\|_{\mathbb{T}} \leq \delta_i \quad (n \in Q), \quad (1.3)$$

where d_Q is the actual positive integer step of Q . The elementary theorem is driven by a checkable simultaneous quadratic recurrence certificate. Combining it with [73, Theorem 1.1] gives, for every fixed q ,

$$|Q| \in [R^\gamma, 2\lceil R^\gamma \rceil - 1], \quad \delta_i = R^{-\gamma}, \quad 0 < \gamma < \frac{1}{(q+1)(42q+1)}, \quad (1.4)$$

for all sufficiently large R , uniformly in the coefficients. The threshold is not numerically explicit, because the cited recurrence theorem does not provide a numerical threshold. The certificate version has no such hidden existence condition.

A probabilistic obstruction proves that a universal full-partition statement with the same balanced powers must have $\gamma \leq 1/(3q + 1)$. This produces a well-defined optimization problem rather than an unsupported optimality claim.

F.21.3 Arbitrary-dimensional cube coefficients

For $d \geq 2$ and real centered f on a finite abelian group, let

$$\mathcal{Q}_d(\delta + tf) = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{\epsilon \in \{0,1\}^d} (\delta + tf(x + \epsilon \cdot h)).$$

The first three nonconstant coefficients vanish. The fourth coefficient is determined by

$$P_d = \frac{6^d - 2 \cdot 4^d + 2^d}{8}, \quad T_d = \frac{8^d - 3 \cdot 6^d + 3 \cdot 4^d - 2^d}{24}. \quad (1.5)$$

It equals $\delta^{2^d-4}(P_d Q(f) + T_d Q_2(f))$, where $Q(f) = \|f\|_{U^2}^4$ and Q_2 is its nontrivial order-two Fourier contribution. Section III.39 proves this statement and a fifth-order formula with a further order-three term over odd-order groups.

F.21.4 What the results establish

The claims in this article have three different dependencies.

Result	Dependence	Quantitative status
Window localization	Finite differences and Gowers–Cauchy–Schwarz	Explicit finite constants
Direct step partition	An explicitly supplied integer certificate	Explicit finite algorithm
Exponent (1.4)	The certificate theorem plus Lau’s recurrence theorem	Uniform asymptotic threshold
Cube coefficients	Finite Fourier analysis and binary configurations	Exact identities

The relevant Lean source interfaces have been read, but the new results here have not been checked by a Lean kernel. Section V.65 specifies the mathematical interfaces suitable for formalization. Section V.67 states further questions and identifies which are local optimization problems and which require the full Szemerédi induction.

F.22 Source 17. The target and the main conclusions

Source 17, delivered Section 1.

Added in the fourth write, 6 October 2026. Delivered title block. “Polynomial-loss restriction in Gowers’s Szemerédi argument / Fejér filters, density-sensitive moments, and finite-field systems / Research note prepared for the ProveIt project / 6 October 2026”; the same in the PDF author field; running head “Polynomial-loss restriction / Gowers’s Section 9 and beyond”. Archive `Gowers_Polynomial_Restriction_Research.zip` (arrival 060528502), manuscript `article.tex` (1,327 lines, 21 pages), pin `dca9f0638` (its `source_manifest.json`, shipped as `data/17-polynomial-restriction-source_manifest.json`). Printed below: the abstract, the framed status statement after it, and Section 1; Sections 2–8 form a chapter of Part II (Section II.57), Sections 9–12 and Appendices A–B are in Part V. Its macro `\code` (verbatim file names) is renamed `\codeseventeen`.

Abstract of source 17 (as delivered). We give a quantitative replacement for the random-restriction step in Section 9 of Gowers’s 2001 proof of Szemerédi’s theorem. A single normalized Fejér kernel replaces a long product of low-degree cosine filters. For a balanced relation with $m = 2r \geq 4$ variables, the method retains at least

$$\frac{\alpha^m \eta^{m-1}}{2(m+2)^m} |B|^{m-1}$$

relations on a subset where at least a $1 - \eta$ proportion are respected. The hypotheses include an explicit polynomial ambient-size threshold. For $m = 16$, this replaces the exponent 2^{19} in the printed Lemma 9.3 by exponents 16 in α and 15 in η . The result also allows fractional initial weights and vector-valued maps in sufficiently large characteristic. Keeping the exact second Fourier moment improves the input to Corollary 9.4 from $\gamma^7 \beta^6$ to γ^7 . A separate random-linear-map construction works in fixed characteristic and treats arbitrary translation-invariant systems of linear relations; its bounds depend on the rank of the relations broken by the map. Finally, a moment inequality identifies a limitation of the scalar filter-certificate method, not a limitation of all possible restrictions. All assertions specific to this note are proved below, and a reproducible finite-audit script is supplied.

Scope and verification. This is an unrefereed research manuscript giving local quantitative refinements, not a claimed improvement to the global bounds for Szemerédi’s theorem. Bibliographic priority for the refinements has not been established. The proofs are written out, but have not been independently refereed or compiled in Lean. The finite computations reported in Section V.70 are checks, not substitutes for the proofs.

F.22.1 Where the improvement occurs

The target is Gowers’s restriction of a partially additive map to a domain on which almost every relevant additive relation is respected. In the notation of his Lemma 9.3, a set $B \subseteq \mathbb{Z}_N$ has size βN , and at least $\alpha \beta^{15} N^{15}$ additive sixteen-tuples are respected by a map ϕ . The printed conclusion retains $(\alpha \eta / 4)^{2^{19}} \beta^{15} N^{15}$ additive sixteen-tuples while making their respected proportion at least $1 - \eta$ [1, Section 9, pp. 515–518]. We work in the prime cyclic setting used in that argument and then state precisely which extensions are valid.

There are two independent losses to address. The first is the extremely small distinguishing advantage of a degree-one cosine filter on a sixteen-variable relation. The second is the replacement of $|B|$ by N in a second-moment estimate before interpolation. The former loss is removed by a higher-degree filter; the latter by keeping the actual mass throughout the calculation.

The resulting prime-cyclic statement is the following. Its full weighted version, and its proof, appear in Theorem II.61.2.

Theorem F.22.1 (Prime-cyclic restriction, headline form). *Let $m = 2r \geq 4$, let N be prime, let $B \subseteq \mathbb{F}_N$ have size $\beta N > 0$, and let $\phi : B \rightarrow \mathbb{F}_N$. Suppose that $0 < \alpha, \eta \leq 1$ and that at least $\alpha|B|^{m-1}$ ordered tuples satisfy both*

$$x_1 + \cdots + x_r = x_{r+1} + \cdots + x_m, \quad \phi(x_1) + \cdots + \phi(x_r) = \phi(x_{r+1}) + \cdots + \phi(x_m).$$

Set

$$L = \left\lceil \frac{m+1}{\alpha\eta} \right\rceil. \quad (\text{F.22.1})$$

If

$$N \geq 2L^m(2L-1)^{m-1}\beta^{-(m-1)}, \quad (\text{F.22.2})$$

there is a subset $B' \subseteq B$ on which at least a $1 - \eta$ proportion of the additive m -tuples are respected and which contains at least

$$\frac{|B|^{m-1}}{2\eta L^m} \geq \frac{\alpha^m \eta^{m-1}}{2(m+2)^m} |B|^{m-1} \quad (\text{F.22.3})$$

additive m -tuples. In particular,

$$|B'| \geq \left(\frac{\alpha^m \eta^{m-1}}{2(m+2)^m} \right)^{1/(m-1)} |B|. \quad (\text{F.22.4})$$

The characteristic condition needed later is automatic here: the right side of (F.22.2) is larger than $2L - 2$. A simpler, larger sufficient threshold is

$$N \geq 2^m(m+2)^{2m-1}(\alpha\eta)^{-(2m-1)}\beta^{-(m-1)}. \quad (\text{F.22.5})$$

Thus the phrase “sufficiently large” has an explicit polynomial meaning for each fixed number of variables.

F.22.2 The order-eight consequence

Putting $m = 16$ gives

$$T_{16}(B') \geq \frac{\alpha^{16}\eta^{15}}{2 \cdot 18^{16}} \beta^{15} N^{15}. \quad (\text{F.22.6})$$

Here and below T_m counts all ordered additive m -tuples, including repetitions; the proof does not silently discard them from the final conclusion. The comparison is with the printed local restriction bound, not with the current global quantitative theory.

The density-sensitive moment calculation proves that if the number of respected additive quadruples is at least $\gamma|B|^3$, then the number of respected additive sixteen-tuples is at least $\gamma^7|B|^{15}$. Consequently (F.22.6) can be applied with $\alpha = \gamma^7$:

$$T_{16}(B') \geq \frac{\gamma^{112}\eta^{15}}{2 \cdot 18^{16}} \beta^{15} N^{15}. \quad (\text{F.22.7})$$

No extra factor β^{96} is introduced by the interpolation input. The corresponding convenient sufficient threshold is

$$N \geq 2^{16} 18^{31} \gamma^{-217} \eta^{-31} \beta^{-15}. \quad (\text{F.22.8})$$

F.22.3 What is, and is not, being claimed

The Fejér kernel, Fourier orthogonality, interpolation of moments, and random linear maps are standard tools. The contribution here is the specific restriction estimate, its complete exceptional-set accounting, the weighted and vector-valued extension, the rank-sensitive system version, and the identification of the corresponding proof-method limitation. The note does not claim that these tools themselves are new.

The high-dimensional fixed-characteristic results in Section II.63 are different from the prime-cyclic theorem. One cannot apply a high-dimensional rank argument to the one-dimensional space $\mathbb{F}_N$ when there are many variables. Conversely, one cannot use a high-degree trigonometric filter in small characteristic without checking Fourier aliasing.

For perspective, Leng, Sah, and Sawhney prove the global bound $r_k(N) \ll N \exp(-(\log \log N)^{c_k})$ for $k \geq 5$, by a different combination of inverse theorems and density increments [2, 3]. No assertion in this note improves that bound. Gowers also explicitly presents Section 9 as a preparatory special case for more complicated later configurations [1, p. 515]; improving it alone does not propagate automatically through the final induction.

[write] **6 October 2026; sources 15–20.** *Formal status, which the source does not state* (checked at HEAD 0b4890a1c). `lemma_9_3` (Sections08_09.lean, line 77) quantifies over all N with `[NeZero N]`, without primality, and has the exact companion `lemma_9_3_holds` (Proofs09Restriction.lean, line 376), proved through `lemma_15_5_zero_dimensional_holds` (Proofs15Restriction.lean, line 3050) with the retention $(\alpha\eta/4)^{2^{19}}\beta^{15}N^{15}$; `corollary_9_4` (line 91) has the exact companion `corollary_9_4_holds` (line 416). Both were exact at the pin. The definitions match this source exactly: T_{16} is `additiveTupleCount 8` and G_{16} is `phiAdditiveTupleCount 8` (Definitions.lean, lines 322 and 327: ordered 16-tuples with repetitions and signs $(1^8, -1^8)$), and “ $G \geq (1 - \eta)T$ ” is `GammaHomOfOrder 8 B' phi (1 - eta)` (line 334). So Theorem F.22.1 at $m = 16$ proves the case N prime of `lemma_9_3` with the explicit threshold $N_0 = \lceil 2L^{16}(2L - 1)^{15}\beta^{-15} \rceil$, $L = \lceil 17/(\alpha\eta) \rceil$: the hypothesis forces $\alpha, \beta \leq 1$, and then $\alpha^{16}\eta^{15}/(2 \cdot 18^{16}) \geq (\alpha\eta)^{16}/(2 \cdot 18^{16}) \geq (\alpha\eta/4)^{2^{19}}$. Likewise (F.22.7) gives the prime case of `corollary_9_4`. It does not prove the all- N propositions, which the development already proves with the weaker bound (Section V.69 says so). Nothing of this source is formalized. *Credit:* source 04’s Sections II.8–II.9 (in the repository at this source’s pin) prove the moment lifting and a polynomial restriction theorem over prime fields; the comparison is in the notes after Proposition II.59.1 and Theorem II.61.2.

F.23 Source 18. Results, scope, and their relation to the source

Source 18, delivered Section 1.

Added in the fourth write, 6 October 2026. Delivered title block. “Higher-Order Cube Expansions and Sharp Degeneracy Bounds in Gowers’ Szemerédi Framework / Simultaneous polynomial partitions, exact constants, and quantitative obstructions / Research report prepared for Vladimir Reshetnikov / 6 October 2026”; the same in the PDF author field. Archive `gowers_cubes_degeneracy_partitions.zip` (arrival 82bc38a30), manuscript `gowers_cubes_degeneracy_partitions.tex` (2,672 lines, 35 pages), pin `dca9f0638` (its `provenance.json`, shipped as `data/18-arrangement-degeneracy-provenance.json`). Printed below: the abstract, its keywords and Section 1; Section 2 is in Part II (with source 19, Section II.65), Section 3 in Part III, Sections 4–5 in Part I, Sections 6–9 and Appendix A in Part V. Its macro `\T` ($\mathbb{R}/\mathbb{Z}$) is renamed `\Tcirc` (the report’s $\mathbb{T}$ is the same circle written $\mathbb{T}$).

Abstract of source 18 (as delivered). We develop three explicit refinements of the local machinery in Gowers’ *A new proof of Szemerédi’s theorem*. First, for the full definition of degeneracy in Section 15, we determine the sharp first-order density of degenerate d -arrangements.

For fixed $k \geq 1$, $d \geq 2$ and primes $p \rightarrow \infty$, this density is $(k + C_d)/p + O_{k,d}(p^{-2})$, where $C_d = \frac{1}{2}(T_{2d} - 2\binom{2d}{d} + 1)$ and T_n is a central trinomial coefficient. In particular $C_8 = 2\,585\,444$. The proof uses paired horizontal and vertical moment equations and gives explicit finite bounds and a direct improvement of the ambient-size hypothesis in Lemma 15.5. Second, we compute the complete fourth- and fifth-order expansion of real cube moments in every dimension, over every finite abelian group. The formula identifies all corrections from characters of orders two and three, and yields a sharp fourth-order leading bound, including sharpness for indicator sets. Third, we give an elementary transfer from simultaneous monomial recurrence to partitions on which several polynomial phases have small oscillation. Existing recurrence theorems yield a polynomial dependence on the number of phases in place of the exponential dependence displayed in Lemma 5.9; a shared-leading-coefficient quadratic family admits a stronger source-internal specialization. An independent volume argument proves a necessary exponent frontier for every polynomial degree. Complete proofs, exact finite checks, a source-level Lean integration plan, and further research questions are included. These are local mathematical results; the report does not claim a new global Szemerédi bound or priority over the entire literature.

Keywords. Gowers norms; Boolean cube circuits; additive arrangements; polynomial recurrence; finite-field rank; Szemerédi’s theorem.

The aim is to improve particular statements in Gowers’ proof while retaining enough information to say exactly what has been gained. The three parts of this report address different losses: the treatment of exceptional arrangements in Section 15, the triangle inequality for cube counts in Section 3, and the repeated refinement of polynomial partitions in Section 5. Each produces a theorem with a full proof, and each has a different consequence for the surrounding argument.

Throughout, references such as “Lemma 15.4” refer to Gowers’ paper [1], not to the numbering of the present article. The repository supplied with the research request contains both a reconstructed $\LaTeX$ source and a substantial Lean development [88]. We read the relevant source files and compared the formulations with the published paper. The reconstructed source includes corrections and explicit hypotheses absent from some printed statements; these are relevant when a local improvement is substituted into a later lemma.

F.23.1 A sharp asymptotic for degenerate arrangements

A d -arrangement in $\mathbb{F}_p^{k+1}$ consists of $2d$ labelled k -cubes with common side vector, whose last coordinates satisfy an additive relation. It has

$$D = (2d + 1)k + 2d - 1$$

free scalar parameters. There is a universal alternating dependence among its vertices. Degeneracy means that some other ternary dependence annihilates every coordinate monomial in the definition. Gowers’ Lemma 15.4 bounds the number of degenerate arrangements by

$$k 3^{2d} 2^k p^{D-1}. \tag{F.23.1}$$

The full definition, rather than the weaker list of tests mentioned after the lemma, permits the following substantially more precise conclusion:

$$\#\{\text{degenerate } d\text{-arrangements}\} = (k + C_d)p^{D-1} + O_{k,d}(p^{D-2}), \quad C_d = \frac{T_{2d} - 2\binom{2d}{d} + 1}{2}, \tag{F.23.2}$$

where

$$T_n = [z^0](z^{-1} + 1 + z)^n.$$

The coefficient $k + C_d$ is optimal. Its k part counts zero-side hyperplanes, and its C_d part counts distinct additional relations among the last coordinates. Every other kind of certificate contributes only at second order or smaller. This separation is what the general polynomial zero estimate in the original proof does not record.

For the 8-arrangements used later in the proof, put

$$M = 2^{k+4}, \quad E_k = 2^{M+k+3}, \quad \lambda = (\alpha\eta/4)^{E_k} \beta^{15}.$$

In the usual range $0 < \alpha, \beta, \eta \leq 1$, a sufficient condition for the degeneracy absorption in Lemma 15.5 is improved from

$$\lambda p \geq k3^M \quad \text{to} \quad p \geq 17, \quad \lambda p \geq k + 2585445. \quad (\text{F.23.3})$$

The very large sampling exponent E_k remains the same. Thus the report improves a specific finite threshold, and identifies its asymptotically correct first-order coefficient, without attributing a further consequence to parts of the proof that have not been re-estimated.

F.23.2 Exact cube coefficients in every dimension

For a real function g on a finite abelian group G , let

$$\mathcal{C}_\ell(g) = \mathbb{E}_{x, h_1, \dots, h_\ell \in G} \prod_{v \in \{0,1\}^\ell} g(x + v \cdot h) = \|g\|_{U^\ell}^{2^\ell}.$$

We use ℓ here to avoid confusing cube dimension with the arrangement parameter d . If $\mathbb{E}f = 0$ and $m = 2^\ell$, the first three nonconstant coefficients in $\mathcal{C}_\ell(\delta + tf)$ vanish. The fourth coefficient is an explicit combination of the ordinary Fourier fourth moment and its restriction to characters of order two. The fifth is a combination of a five-circuit Fourier correlation and an order-three correction. The four counting constants are

$$\begin{aligned} R_\ell &= \frac{6^\ell - 2 \cdot 4^\ell + 2^\ell}{8}, & B_\ell &= \frac{2^\ell(2^\ell - 1)(2^\ell - 2)}{24}, \\ Q_\ell &= \frac{8^\ell - 3 \cdot 6^\ell + 3 \cdot 4^\ell - 2^\ell}{6}, & S_\ell &= \frac{10^\ell - 4 \cdot 8^\ell + 6 \cdot 6^\ell - 4 \cdot 4^\ell + 2^\ell}{24}. \end{aligned}$$

These formulas count the relevant four- and five-vertex configurations in the Boolean cube. Smith normal form explains why only orders two and three can enter at these orders. The result is valid for every finite abelian group; excluding small torsion is a corollary, rather than a tacit assumption in the proof.

For an indicator $\mathbf{1}_A = \delta + f_A$ with $\|f_A\|_{U^\ell} \leq \varepsilon$, the resulting estimate begins at $B_\ell \delta^{m-4} \varepsilon^4$ instead of the linear term in $(\delta + \varepsilon)^m$. For odd-order groups the fourth coefficient can be bounded using R_ℓ instead. The all-group coefficient B_ℓ is sharp, even along sequences of indicator sets with density tending to any fixed $\delta \in (0, 1)$. The proof first gives an exact order-two-character example for real weights, and then transfers it to sets by independent rounding with a quantified Gowers-norm error.

F.23.3 Partitions and their quantitative limits

The polynomial part keeps the near-prescribed lengths of Lemma 5.9: every cell has length $v - 1$ or v . It replaces repeated induction on the number q of phases by an induction on polynomial degree which treats all q leading coefficients together. A recurrence-radius function is the only

analytic input, so the elementary partition argument can be combined with either the source’s scalar estimate or later simultaneous recurrence theorems [72, 73].

For arbitrary quadratic phases, Lau’s stated recurrence estimate gives every strict exponent

$$a < \frac{1}{126q^2 + 129q + 2} \quad (\text{F.23.4})$$

for lengths $v - 1, v$ with $v \leq r^a$ and containing-arc diameter at most r^{-2a} , once r is sufficiently large. The displayed quadratic case of Gowers’ lemma uses $a = 1/(2 \cdot 2048^q)$. The new conclusion has an unspecified coefficient-independent threshold inherited from the cited recurrence theorem. Its numerical exponent should not be confused with a small or fully numerical threshold for r .

For quadratics whose leading coefficients are bounded integer multiples of one common frequency, the scalar estimate in Gowers’ own Lemma 5.5 instead gives every $a < 1/(99q + 98)$, with an explicit, though enormous, threshold. Arbitrary lower coefficients are permitted in this specialization.

There is also a rigorous obstruction. If a universal degree- k , q -phase partition theorem guaranteed all cell lengths $\gg r^a$ and all oscillations $\ll r^{-b}$, with $b > 0$, it would necessarily satisfy

$$\left(\frac{qk(k+1)}{2} + 1 \right) a + kqb \leq 1. \quad (\text{F.23.5})$$

For quadratics and $b = 2a$, this says $a \leq 1/(7q + 1)$. The sufficient q^{-2} rate and this q^{-1} obstruction leave a substantive research gap. No optimality assertion for the general quadratic partition exponent is made.

F.23.4 What is established, imported, and left open

The rank arguments, coefficient counts, cube expansion, rounding sharpness, recurrence-to-partition transfer, and polynomial obstruction are proved in this article. Gowers–Cauchy–Schwarz, Fourier orthogonality, Smith normal form, and the cited recurrence estimates are established inputs. The familiar twenty circuits of the three-dimensional cube were explicitly described in earlier literature, for example Beerenwinkel–Pachter–Sturmfels [74]; the low-dimensional configuration itself is not presented as a discovery of this report.

The article extends the earlier local calculations to precise all-dimensional statements and pursues the Section 15 and nonlinear partition questions further. The statements can be assessed from their proofs independently of a claim of historical priority. A targeted literature check does not establish that no equivalent result appears elsewhere.

Nor does a collection of local improvements establish a new best bound for Szemerédi’s theorem. For context, Leng–Sah–Sawhney proved, for each $k \geq 5$, a bound of the form

$$r_k(N) \ll_k N \exp(-(\log \log N)^{c_k}), \quad c_k > 0,$$

using quantitative inverse theorems and a density-increment strategy [2]. The results here have a different, explicitly local purpose. Their value includes sharp information about exceptional sets, exact spectral expansions, and more effective interchangeable components for proofs and formalizations.

F.23.5 Conventions

Finite averages are normalized. A finite group has at least one element. Characters are written additively in the dual group; $j\chi = 0$ means that the character’s order divides j . A cube is a labelled

affine map, so coincident vertices are included in cube counts. In the arrangement part, $d \geq 2$ is the number of cubes on either side of the additive relation, and $k \geq 1$ is their horizontal dimension. In the partition part, k is polynomial degree and q is the number of phases. These parameters are reintroduced at the start of each part. Polynomial oscillation always means the length of a containing circle arc, as in Gowers’ diameter convention.

F.24 Source 19. Scope and principal results

Source 19, delivered Section 1.

Added in the fourth write, 6 October 2026. Delivered title page. “Research article / Additive combinatorics / Efficient Interpolation in Gowers’ Szemerédi Framework / Exact sampling losses, shared refinements, and simultaneous polynomial recurrence / Prepared for Vladimir Reshetnikov / 6 October 2026”; *the PDF author field reads “ChatGPT, research report prepared for Vladimir Reshetnikov”*. Archive `gowers_interpolation_article.zip` (arrival 82bc38a30), manuscript `gowers_interpolation.tex` (2,505 lines, 32 pages), pin `dca9f0638` (its `provenance/source-audit.json`, shipped as `data/19-shared-interpolation-source-audit.json`). Printed below: the abstract, the closing statement of its title page and Section 1; Sections 2–5 and 8 are in Part II (Section II.65), Section 6 in Part I, Sections 7, 9 and 10 in Part V. Its ‘question’ environment (printed “Research question”) is set in this report’s ‘Research question’ style; its macro `\commit` (its own pin) is renamed `\commitnineteen`; its `cleveref` references are written out as `cleveref` prints them.

Abstract of source 19 (as delivered). We strengthen the interpolation mechanism in Section 16 of Gowers’ proof of Szemerédi’s theorem and give a complementary refinement of its polynomial phase partition lemma. A common sample of $\min\{n, \lceil qd/\sigma \rceil - 1\}$ distinct columns recovers all but σ of the ambient mass in any family of partial rows containing at most q polynomial classes of degree at most $d - 1$. We determine the exact finite minimax loss by a hypergeometric optimization and prove its sharp sparse-sampling asymptotic. For affine interpolation the constant is $\varphi^3 e^{-\varphi} = 0.839962094657\dots$, where $\varphi = (1 + \sqrt{5})/2$. Refining the sampled cross-sections once, before forming their interpolants, replaces a quadratic number of structural refinement steps by a linear number and reduces the list count to an elementary symmetric polynomial in the slice list sizes. The result extends to arbitrary fixed degree in the final variable and includes a proper-box synchronization argument. Separately, Maynard’s established simultaneous recurrence theorem yields polynomial rather than exponential dependence on the number of polynomial phases, for fixed degree. A self-contained volume obstruction gives a necessary degree-sensitive frontier for any such partition theorem. Full written proofs, exact finite diagnostics, a source audit, and a research agenda are supplied. These are local refinements; neither a new global Szemerédi bound nor a new Lean kernel certification is asserted.

The source comparison is pinned to ProveIt commit

`dca9f06383ccdfc2ed208ccb7c9882214fdb09ac`.

The original 2001 publication and the repository’s edited transcription are distinguished throughout. The article is a research report; historical priority for its exact formulations is not claimed.

The difficult part of the inverse machinery in Gowers’ argument is not merely finding many local affine pieces. Those pieces must be related across different fibres, and the resulting descriptions must remain economical after repeated partitioning. Section 16 uses sampled cross-sections and interpolation to make that transition. This is the main target of the present article [1, 89].

Two losses can be treated separately. First, the sampling argument discards small classes and then insists on a small failure probability for every remaining class. Averaging the *mass* of failed

classes gives a better estimate. Second, once interpolation candidates have been formed, treating them as independent functions causes the same sampled data to be refined repeatedly. A common refinement of the sampled cross-sections allows all interpolants to be formed afterward.

These changes are independent of the more difficult global inverse theorem. They give explicit local statements with transparent inputs, and they also suggest exact extremal questions. We solve the finite sampling optimization for the specified interpolation recovery rule and identify its sharp asymptotic constant. A separate part of the article uses modern recurrence theory to improve the dependence on the number of phases in Lemma 5.9.

F.24.1 The main conclusions at a glance

A. Exact finite sampling loss. For an n -point column set, at most q disjoint classes per row, and interpolation order d , define

$$W_{n,r,d}(s) = \frac{s}{n \binom{n}{r}} \sum_{j=0}^{\min(d,s)-1} \binom{s}{j} \binom{n-s}{r-j}.$$

The worst possible loss of the best common r -column sample is exactly

$$\mathcal{M}_{n,q,d}(r) = \max_{\substack{s_1, \dots, s_q \in \mathbb{Z}_{\geq 0} \\ s_1 + \dots + s_q \leq n}} \sum_{i=1}^q W_{n,r,d}(s_i).$$

The loss is at most $\min\{1, qd/(r+1)\}$ and is zero for $r = n$. The equality statement is a genuine finite minimax theorem for the specified common-anchor recovery mechanism.

B. Sharp asymptotic constant. For fixed d , as $r \rightarrow \infty$, $r/n \rightarrow 0$, and $q/r \rightarrow 0$,

$$\mathcal{M}_{n,q,d}(r) \sim \frac{qC_d}{r}, \quad C_d = \max_{\lambda \geq 0} \lambda e^{-\lambda} \sum_{j=0}^{d-1} \frac{\lambda^j}{j!}.$$

In the affine case, $d = 2$, this is $C_2 = \varphi^3 e^{-\varphi}$. The order q/σ is therefore optimal within this sampling model in the sparse-sampling limit.

C. Shared interpolation refinement. If the r sampled cross-sections have list sizes $M_1, \dots, M_r$ on a common partition, then all degree-at-most- $d-1$ interpolants require at most

$$\sum_{j=1}^{\min(d,r)} e_j(M_1, \dots, M_r)$$

candidate graphs, where e_j is the j th elementary symmetric polynomial. With a common list size M , the affine count is $rM + \binom{r}{2}M^2$. The partition is refined once per sampled cross-section, rather than once for every interpolation pair.

D. A degree-sensitive recurrence transfer. Let q_j count the phases whose degree is at least j , and let κ_j be a valid simultaneous recurrence constant of degree j , with $\kappa_1 = 1$. Define $F_0 = 1$, $G_0 = 0$, and

$$F_j = (1 + j\kappa_j q_j)F_{j-1}, \quad G_j = (1 + j\kappa_j q_j)G_{j-1} + \kappa_j q_j.$$

For $F_k a + G_k b < 1$, every sufficiently long integer interval can be partitioned into progressions of length at least N^a on which all phase diameters are at most N^{-b} . Maynard’s published theorem supplies the constants; its recurrence assertion is an external input [72]. For fixed degree, $F_k + G_k = O_k(q^k)$ when $q_j = q$.

E. A necessary polynomial frontier. If the prescribed phase degrees are $d_1, \dots, d_q$, set $D = \sum_i d_i$ and $S = \sum_i d_i(d_i + 1)/2$. A universal progression partition with length exponent a and diameter exponent b must satisfy

$$(S + 1)a + Db \leq 1.$$

The proof uses finite differences and a volume estimate, not an unproved Diophantine approximation conjecture. For nonlinear phases the gap between this obstruction and the sufficient frontier remains open in the present approach.

F.24.2 How the results differ from earlier project refinements

This article concentrates on the conversion of fibrewise polynomial descriptions into compatible descriptions in all variables, and on simultaneous recurrence for several nonlinear phases. It does not repackage the project’s earlier cube expansions, energy inequalities, Bohr extension estimates, or sharp linear phase-density selection results. The relevant source targets are Lemmas 16.9–16.10 and the nonlinear part of Lemma 5.9. The original and edited source statements, and the live Lean interfaces, are examined in section V.80.

[write] **6 October 2026; sources 15–20.** *What this “does not repackage”.* At the pin the report directory held source 04’s manuscript as `article.tex`; sources 05 and 13 (linear phase partitions) were not visible to this source. Within this report, Section I.34 meets source 18’s Sections 4–5 (the same theorem, proved independently: note after Corollary I.32.4) and extends source 13’s linear results; Sections II.67–II.70 are new ground: no other source treats Section 16 of the paper.

F.24.3 Novelty, attribution, and the scope of the conclusions

The tools used here have substantial precedents. Missing-mass optimization and occupancy estimates are established subjects [75, 76]. Lagrange interpolation, finite averaging, and common refinement of partitions are classical. Maynard’s recurrence theorem is imported with attribution; the optional numerical specialization from Lau is explicitly identified as a preprint input [72, 73].

The contributions claimed here are the precise finite minimax formulation, its interpolation-compatible endpoint, the sharp asymptotic analysis for that model, the more economical lifting theorem, and the proved recurrence transfer and obstruction. These results improve the specified local procedures. A complete priority search would be needed before claiming that every individual formulation is new to the literature. No well-known open global conjecture is claimed to be resolved.

For perspective, substantially stronger global bounds than the 2001 estimate are already available. Leng, Sah, and Sawhney prove, for each $k \geq 5$, a bound of the form $r_k(N) \ll N \exp(-(\log \log N)^{c_k})$, where $c_k > 0$ and $r_k(N)$ is the largest size of a k -term-progression-free subset of $[N]$ [2]. The local comparisons here do not improve that result. More fundamentally, the explicit closing graph-budget comparison in the supplied Section 16 proof requires repair; reducing a sampling constant does not discharge it.

F.25 Source 20. The target and the main conclusions

Source 20, delivered Section 1.

Added in the fourth write, 6 October 2026. Delivered title page. “Additive combinatorics / Quadratic Counting Errors and Sharp Directional Bounds in Gowers’s Szemerédi Argument / Two-function uniformity, torsion-sensitive constants, and a sharp three-term obstruction / Prepared for Vladimir Reshetnikov / 6 October 2026”; the PDF author field reads “Prepared for Vladimir Reshetnikov”. Archive `gowers_quadratic_counting.zip` (arrival 83f3171c1), manuscript `article.tex` (1,535 lines, 21 pages); no commit pin: its Appendix A and `source_manifest.json` (shipped as `data/20-quadratic-counting-source_manifest.json`) record the observed branch checkpoint 1bd57e91c and blob identifiers. Printed below: the abstract, the closing statement of its title page and Section 1; Sections 2–8 form a chapter of Part III (Section III.41), Sections 9–11 and Appendix A are in Part V. Its macro `\U` (a Gowers norm) is renamed `\Usemnorm` (the report’s $\mathcal{U}$ is source 02’s functional); its ‘question’ environment (“Question”, numbered on its own counter) is set in this report’s ‘Question’ style, numbered within its section.

Abstract of source 20 (as delivered). We refine the counting stage of Gowers’s proof by retaining two functions through the Cauchy–Schwarz argument. For surjective homomorphisms from a finite abelian parameter group, the resulting bound is a product of two directional Gowers seminorms and one ordinary L^2 norm. A sharp finite-index comparison converts the directional norms to global norms, with explicit constants that record torsion rather than hiding it behind a prime-modulus hypothesis. A second-order telescoping identity then gives quadratic, rather than linear, error in the balanced uniformity norm. For a set of density δ and admissible k -term progressions,

$$|\Lambda_k(1_A) - \delta^k| \leq B_k(\delta) \|1_A - \delta\|_{U^{k-1}}^2, \quad B_k(\delta) \leq \frac{\sqrt{\delta}}{(1-\delta)^2}.$$

We give a sharper explicit formula for B_k , unequal-density versions, and supersaturation consequences. An exact quadratic-phase construction, followed by probabilistic rounding, proves that the exponent two cannot be increased for three-term counts, even for indicator sets on growing $\mathbb{F}_{29}$ -vector spaces. The article includes proofs, limitations, 10,389 exact finite checks, a formalization plan, and twelve research questions. The improvements are relative to specified local estimates; worldwide priority and a new best global Szemerédi bound are not asserted.

Research article and reproducibility package. The supplied proofs are human-readable mathematical proofs; no new Lean kernel verification is claimed.

The target is the passage from uniformity to counting in Section 3 of Gowers’s paper [1]. In normalized notation, its Theorem 3.2 controls a progression average by the U^{k-1} norm of one function. Corollary 3.3 expands balanced functions and obtains a counting estimate with a factor 2^k . Its hypotheses are asymmetric: the first two sets need not be uniform. The ProveIt transcription makes the coefficient invertibility conditions explicit [94].

The inspected research article in the repository already improves this expansion by density-weighted, first-order telescoping [93]. For a single set A of density δ , that argument gives

$$|\Lambda_k(1_A) - \delta^k| \leq S_{k-2}(\delta)u, \quad u = \|1_A - \delta\|_{U^{k-1}}, \quad S_{k-2}(\delta) = \sum_{r=0}^{k-3} \delta^r. \quad (\text{F.25.1})$$

The improvement developed here is not another reordering of this linear estimate. We retain a second uniformity factor before estimating the balanced expansion.

[write] 6 October 2026; sources 15–20. *Framing within this report.* The “inspected research article” is source 04’s delivered manuscript (blob 45b408f, staged as `article.tex` by aefb0b44a and inspected at the checkpoint 1bd57e91c, three minutes before the first write replaced it). Against that text the sentence is accurate. But the quadratic estimate is not new in this report: source 06’s Corollary III.5.5 proves $|\Lambda - \delta^k| \leq \sqrt{\delta} B_k^{(06)}(\delta) \varepsilon^2$ with the same norm, and its Proposition III.5.6 the sharp three-term coefficient $\sqrt{\delta(1-\delta)}$, which is this source’s B_3 . Source 06 was placed (ad37962c4) as a merge member whose manuscript was not staged, so this source could not see it; the two are independent. New here: a coefficient strictly smaller than source 06’s for every $k \geq 3$ and $0 < \delta < 1$ (note after Corollary III.45.4), the generality of surjective homomorphisms with torsion constants, the sharp finite-index comparison, the pair correction, and quadratic sharpness at fixed density for indicator sets.

F.25.1 A concrete statement to keep in view

Let G be a finite abelian group, $|G| = N$, and let

$$\Lambda_k(a) = \mathbb{E}_{x,y \in G} \prod_{r=0}^{k-1} a(x + ry).$$

Call the progression system *admissible* when multiplication by every integer $1, \dots, k-1$ is an automorphism of G . In a cyclic group this means $\gcd(N, (k-1)!) = 1$, not necessarily that N is prime.

For $0 < \delta < 1$, write $\sigma = \sqrt{\delta(1-\delta)}$ and define

$$R_m(\delta) = \min \left\{ \sqrt{\delta}, \sigma \frac{1 - \delta^m}{1 - \delta} \right\}, \quad m \geq 1, \quad (\text{F.25.2})$$

$$B_k(\delta) = \sum_{j=2}^{k-1} (j-1) \delta^{j-2} R_{k-j}(\delta). \quad (\text{F.25.3})$$

We prove

$$\boxed{|\Lambda_k(1_A) - \delta^k| \leq B_k(\delta) u^2.} \quad (\text{F.25.4})$$

The coefficient is explicit, and satisfies

$$B_3(\delta) = \sqrt{\delta(1-\delta)}, \quad B_k(\delta) \leq \frac{\sqrt{\delta}}{(1-\delta)^2}. \quad (\text{F.25.5})$$

For $0 < \delta \leq 1/2$ there is the simpler exact expression

$$B_k(\delta) = \sqrt{\delta} \sum_{r=0}^{k-4} (r+1) \delta^r + (k-2) \delta^{k-3} \sqrt{\delta(1-\delta)}. \quad (\text{F.25.6})$$

An empty sum, here and below, is zero. The cases $\delta = 0, 1$ have zero error and are treated separately.

The two estimates (F.25.1) and (F.25.4) can be used together by taking their minimum. The quadratic estimate is particularly useful when u is small. It requires uniformity at two marked positions in each surviving term. Consequently it is a strengthening for the symmetric counting problem, not a replacement of the asymmetric Corollary 3.3 under unchanged hypotheses.

F.25.2 What is proved, and what remains distinct

The principal structural result is Theorem III.43.3, for arbitrary surjective homomorphisms of finite abelian groups. Theorem III.44.1 determines the optimal constant in a directional-to-global norm comparison for every fixed tuple of subgroups. Theorem III.45.2 and Corollary III.45.4 give the counting refinements. Theorem III.48.3 establishes sharpness of the exponent two for actual three-term counts on a specified family of groups.

These statements resolve several local questions posed by the argument: linear dependence on the balanced norm is not intrinsic to symmetric counting; torsion can be included quantitatively; the subgroup norm conversion has an exact sharp constant; and, for three-term counts, quadratic dependence cannot be replaced uniformly by any higher power. The last assertion is about finite abelian groups in a growing-dimensional family. We do not claim the corresponding sharpness construction here for prime cyclic groups, or for $k \geq 4$.

The underlying tools—Cauchy–Schwarz, Fourier orthogonality, and quadratic Gauss sums—are classical. Directional Gowers norms have an established theory, including the substantially deeper concatenation results of Tao and Ziegler [77]. Our comparison concerns finite-index subgroups and needs no concatenation theorem. A targeted source check does not establish worldwide priority for the exact formulations below; they should be viewed as proved refinements relative to the specified source estimates.

Nor does a sharper counting lemma by itself establish a new record for $r_k(N)$. For perspective, Leng, Sah, and Sawhney prove $r_k(N) \ll N \exp(-(\log \log N)^{c_k})$ for $k \geq 5$, using quasipolynomial inverse bounds and a density-increment strategy [2]. We make no global bound claim beyond the local implications proved in Section III.47.

F.26 Source 21. The question and the contribution

Source 21, delivered Section 1.

Added in the fifth write, 6 October 2026. Delivered title page. “Additive combinatorics / 6 October 2026 / Sharp Characteristic-Five Cube Bounds in Gowers’s Framework / Optimal sixth-order correction, phase rigidity, and an eighth-order extremal term / Prepared for Vladimir Reshetnikov / Mathematical research report prepared by ChatGPT”; *the PDF author field reads “ChatGPT; prepared for Vladimir Reshetnikov”*. Archive `gowers_characteristic_five.zip` (arrival 26908b3b6), manuscript `gowers_characteristic_five.tex` (1,775 lines, 24 pages), pin 43306653a (its `data/source_manifest.json`, shipped as `data/21-characteristic-five-source_manifest.json`). Printed below: the abstract, the principal result and closing statement of its title page, and Section 1; Sections 2–9 form a chapter of Part III (Section III.49), Sections 10–11 and Appendices A–B are in Part V. Its macros `\U`, `\Cfive` and `\betafive` are unused and not defined; its ‘question’ environment (“Question”) is set in this report’s ‘Question’ style.

Abstract of source 21 (as delivered). Let G be a nontrivial finite abelian group of exponent five, let $f : G \rightarrow \mathbb{R}$ have mean zero, and put $u = \|f\|_{U^3}$. We prove, with absolute constants $c, C > 0$ and without a pointwise bound on f , that

$$\|\delta + f\|_{U^3}^8 \leq \delta^8 + 6\sqrt{2}\delta^4 u^4 + \frac{2^{3/4}}{3}\delta^2 u^6 + Cu^8 \quad (\delta > 0, u \leq c\delta).$$

The coefficient at u^6 is optimal even on a fixed five-point quotient. This closes the characteristic-five coefficient gap in an earlier project report, which established the exponent and a matching lower candidate but only an unspecified upper constant.

The proof strengthens the treatment of the sixth- and seventh-degree terms in the centered cube expansion. A positive cubic convolution statistic controls both terms, and its Fourier-tail estimate permits their absorption against a retained norm deficit. Four further Fourier squares yield the sharp phase-dependent coefficient

$$\frac{2^{3/4}}{3} \frac{(\cos \theta)_+^2}{1 + \cos^2 \phi}, \quad \theta = 2 \arg \widehat{f}(\xi) - \arg \widehat{f}(2\xi), \quad \phi = 3 \arg \widehat{f}(\xi) + \arg \widehat{f}(2\xi),$$

when $\{\pm\xi\}$ is a largest conjugate Fourier pair. Near equality forces a second harmonic of size $2^{1/4}u^2/(6\delta)$ and the phase condition $e^{10i \arg \widehat{f}(\xi)} \rightarrow -1$.

On $G = \mathbb{Z}/5\mathbb{Z}$, an exact six-row Fourier formula and analytic optimization determine the next term of the weighted supremum:

$$M_{\delta,5}(u) = \delta^8 + 6\sqrt{2}\delta^4u^4 + \frac{2^{3/4}}{3}\delta^2u^6 + \frac{20}{9}u^8 + O_\delta(u^{10}).$$

We transfer sharpness to indicator sets in growing exponent-five groups, give explicit finite profile inequalities, explain the interface with Gowers’s Lemmas 3.8–3.10 and the repository, and distinguish the remaining higher-dimensional eighth-order problem from the questions resolved here.

Principal result. On every nontrivial finite abelian group of exponent five, the sharp next coefficient after $6\sqrt{2}\delta^4u^4$ in the centered three-cube problem is $2^{3/4}\delta^2/3$, at order u^6 . The uniform error is $O(u^8)$. On the five-point group the following coefficient is exactly $20/9$.

Based on the supplied Ramsey development in `VladimirReshetnikov/ProveIt`, inspected at commit `43306653afaaaaa0b58a7ca71e4067c245594867`. Complete written proofs, exact finite formulas, computational checks, and a further research agenda are included. The report solves a specified extremal question left open in an earlier project manuscript; it does not assert bibliographic priority or a new global quantitative Szemerédi bound. The new proofs have not been certified by Lean.

Gowers’s proof of Szemerédi’s theorem develops higher uniformity norms to measure additive structure [1]. Lemma 3.10 bounds the number of d -dimensional additive cubes in a set of density δ by

$$(\delta + \alpha^{1/2^d})^{2^d} N^{d+1},$$

where α is the degree- $(d-1)$ uniformity parameter. Its proof uses the triangle inequality from Lemma 3.9, itself a consequence of the cube Cauchy–Schwarz inequality in Lemma 3.8. This is an effective general estimate, but it does not exploit the real-valued, mean-zero structure of the balanced function.

The Ramsey research collection in the supplied repository [103] already contains extensive local improvements. In particular, the quartic leading term in the real, centered three-cube problem and its Fourier concentration mechanism are prior project results. An additional manuscript, *Sharp Extremal Refinements in Gowers’s Szemerédi Framework* [102], determines a secondary $u^{16/3}$ term over suitable odd-order groups. Its Section 3 shows that exponent five changes the next exponent to 6 and constructs the lower coefficient $2^{3/4}/3$, while explicitly leaving the optimal upper coefficient unresolved.

[write] **6 October 2026; sources 21–27.** *Identifications and credit within this report.* The “research collection” [103] at this source’s pin is the first write of this report (sources 01–05); the quartic leading term and its Fourier concentration are source 04’s Theorem III.3.4 and its consequences. The manuscript *Sharp Extremal Refinements in Gowers’s Szemerédi Framework* [102] is source 14; its Section 3 is Section III.35,

with the exponent six and the lower coefficient in Theorem III.35.2. The coefficient question settled below is source 14's Question 1 (item III.23). Source 23, written the same day from a pin three minutes later, proves the same theorem independently (its Theorem III.58.1); neither cites the other. In this report source 21 is the main source and source 23 a credited second route (Section III.49).

The present report develops that unresolved direction. It does not relabel the existing quartic estimate or the previous lower example as new. Instead, it proves a matching uniform upper coefficient, a sharp phase envelope, a stronger error estimate, and an exact further term on the smallest group. Background inequalities are reproved because their precise constants and equality defects are needed.

F.26.1 Statements at a glance

All averages in this report are probability averages. For real W , write

$$Q_3(W) = \mathbb{E}_{x,h_1,h_2,h_3 \in G} \prod_{\omega \in \{0,1\}^3} W(x + \omega \cdot h) = \|W\|_{U^3}^8.$$

The constants used repeatedly are

$$C_5 = \frac{2^{3/4}}{3} = 0.5605976101\dots, \quad \beta_5 = \frac{2^{1/4}}{6} = 0.1982011858\dots \quad (\text{F.26.1})$$

Theorem F.26.1 (Uniform characteristic-five expansion). *There are absolute constants $c, C > 0$ such that, for every nontrivial finite abelian group G of exponent five, every real mean-zero f , and every $\delta > 0$, putting $u = \|f\|_{U^3}$, one has*

$$Q_3(\delta + f) \leq \delta^8 + 6\sqrt{2}\delta^4 u^4 + C_5 \delta^2 u^6 + C u^8 \quad \text{if } u \leq c\delta. \quad (\text{F.26.2})$$

For every fixed $0 < \delta < 1$ and every fixed nontrivial such group, $[0, 1]$ -valued weights of mean δ attain the displayed coefficients through order u^6 . Both the constants in the upper bound and the smallness condition are independent of the dimension of G .

[write] **6 October 2026; sources 21–27.** *Second route, and a check.* Source 23's Theorem III.58.1 is the same statement, with the supremum over weights on every fixed group and over all such groups; its constants are cruder (its sextic tail term $84t^6\sqrt{D}$ against $42\sqrt{3}t^6D^{1/2}$ here), and its rigidity statement is equivalent to Theorem III.54.1 (note there). The coefficient was re-derived at the write: $\sup_{b \geq 0} \{-12 \cdot 2^{1/4}b^2 + 4\sqrt{2}b\} = (4\sqrt{2})^2/(48 \cdot 2^{1/4}) = 2^{3/4}/3$, attained at $b = 2^{1/4}/6 = \beta_5$. The theorem answers source 14's Question 1, including its sub-question on the rank of the group: the bound is uniform, and the coefficient is attained on every fixed group (Corollary III.55.2).

Theorem III.53.2 refines C_5 to an explicit function of two Fourier phases. Theorem III.54.1 gives necessary conditions for asymptotic equality. Theorem III.55.1 proves the additional coefficient $20/9$ on $\mathbb{Z}/5\mathbb{Z}$, including a proof that the formula describes the global maximum for small u .

Question		Prior project information	Result proved here
Exponent-five	sec-	Order u^6 ; lower coefficient C_5 ; upper coefficient unspecified	Sharp coefficient C_5 and uniform remainder $O(u^8)$
ondary term			
Phase dependence		One phase choice gives a lower example	Exact asymptotic phase envelope and phase rigidity
Fixed five-point	next	Not determined in the inspected	Coefficient $20/9$, remainder $O_\delta(u^{10})$
term		sources	
Indicator sharpness in		Weighted candidate	Indicator families attaining the same
exponent five			coefficients through u^8 as lower examples

F.26.2 Scope and comparison with the literature

The optimization problem is local in the uniformity norm. It asks how large the cube count of a small perturbation of a constant can be when the norm of that perturbation is fixed. Eisner and Tao [7] study a different extremizer problem: Gowers–Host–Kra seminorms that are large under L^∞ or critical L^p constraints. Hatami [31] develops norm geometry and Hanner inequalities for generalizations of Gowers norms. These works provide appropriate context for rigidity and convexity, but their constraints should not be identified with the fixed-mean, small-centered-norm problem here.

The literature search did not establish bibliographic priority for the new formulas. Accordingly, “sharp” always means that the stated inequality has matching examples in its specified mathematical class; it does not mean that no related observation exists elsewhere. The unresolved question settled here is explicitly the coefficient question in the earlier project report.

No new bound for the largest progression-free subset of $\{1, \dots, N\}$ is obtained. As a quantitative benchmark, Leng, Sah, and Sawhney [2] prove for each fixed $k \geq 5$ a bound $r_k(N) \ll_k N \exp(-(\log \log N)^{c_k})$ with $c_k > 0$. The present characteristic-five results concern a different local step and do not propagate through a new global density-increment argument. In particular, a group of exponent five has no progression with more than five distinct points; its cube geometry remains meaningful independently of that restriction.

F.26.3 What is counted

A three-cube is an ordered parameter tuple $(x, h_1, h_2, h_3) \in G^4$. Its vertices may coincide, and a sidelength may be zero. Thus $|G|^4 Q_3(1_A)$ is an integer count of parameterized cubes, not a count of unlabelled eight-point subsets. This convention is the one relevant to the uniformity norm and to Lemma 3.10.

Every nontrivial finite abelian group of exponent five is a finite dimensional vector space over $\mathbb{F}_5$. We use additive notation for both G and its character group $\widehat{G}$; a dual element r is also the character $x \mapsto r(x)$ when it occurs in Fourier inversion. No coordinate basis is needed in the proofs.

F.27 Source 22. Purpose, results, and scope

Source 22, delivered Section 1.

Added in the fifth write, 6 October 2026. Delivered title block. “Degree Filtrations in Gowers’ Ramsey Framework / Sharper degeneracy bounds, exact polynomial pattern thresholds, and obstructions to extending norm constructions / Prepared for Vladimir Reshetnikov / Research report developed with ChatGPT / 6 October 2026”; *the PDF author field reads “ChatGPT; prepared for Vladimir Reshetnikov”*. Archive `gowers_degree_stratification.zip` (arrival 26908b3b6), manuscript `gowers_degree_stratification.tex` (2,071 lines, 30 pages), pin 43306653a (its `SOURCE_LEDGER.json`, shipped as `data/22-deg-ree-filtrations-SOURCE_LEDGER.json`). Printed below: the abstract and Section 1; Sections 2–5 are in Part II (with sources 25 and 27, Section II.81), Sections 6–10 form a chapter of Part III (Section III.60), Sections 11–12 are in Part V. Its ‘question’ environment (“Research question”) is set in this report’s ‘Research question’ style; three of its statements, unlabelled as delivered, carry labels added at the write.

Abstract of source 22 (as delivered). We develop two degree-sensitive refinements connected with Gowers’ proof of Szemerédi’s theorem. First, we stratify the exceptional certificates in Section 15 by the Boolean degree left after removing parity. A degree- t certificate has success probability at most p^{-t-1} for $t \geq 2$. For degree one we count exactly the certificate classes satisfying the necessary balance constraints. Their number is a polynomial in the cube dimension k of degree $2d$, rather than an exponential in 2^k . For $d = 2$ it is $30k^4 + 204k^3 + 111k^2 + 114k$. This yields a fully explicit finite replacement for the exceptional-set estimate, with degree-sensitive sufficient field-size conditions for its absorption step. We also settle the first-order degeneracy count for the previously excluded two-cube case: its optimal coefficient is $k^2 + 3k$.

Second, for any prescribed equality pattern on a finite set of field parameters, we determine the exact minimum number of successful directions at a prescribed center among polynomial colorings of given degrees. The answer is $q^{\max(n-B,0)}$, where B is a maximum-weight column-basis sum of an explicit evaluation matrix. This removes the earlier degree restriction for symmetric progressions and gives a monochromatic analogue. Norm forms attain the pointwise bound. We prove that a complete low-degree norm construction cannot be extended by adding an arbitrary independent summand; in the four-term case we give an exact count in terms of third differences of that summand. Quantitative constant-line counts explain a further obstruction in the characteristic-divisible norm layers. All principal statements have proofs; exact finite verification and integration notes accompany the source. No improved global Szemerédi bound or resolution of the global quintic avoidance problem is asserted.

The useful question behind this report is whether the degree actually present in an auxiliary object can replace the worst possible degree implicit in a large counting argument. We answer this in two settings. In Gowers’ exceptional-arrangement argument, degree controls the number of independent restrictions imposed by a certificate. In polynomial coloring, degree controls the number of variables that can be constrained by the evaluation equations. Both arguments become substantially more precise when they retain a filtration by degree.

Our direct source is Section 15 of Gowers’ 2001 paper [1], as transcribed and studied in the Ramsey part of ProveIt [104]. We use repository snapshot 43306653afaa; the full revision identifier appears in the references and source ledger. The earlier project reports [105, 106] already determined the first-order coefficient for $d \geq 2$ arrangements and the low-degree symmetric-coloring threshold. Those results are prerequisites and comparison points, not claimed again as new discoveries. The necessary arguments are reproduced here so that this article can be read alone.

[write] **6 October 2026; sources 21–27.** *Identifications and credit within this report.* [105] is source 18: its Theorem II.66.8 (the coefficient $k + C_d$), Proposition II.66.3 and Lemma II.66.2 (the paired rank argument). [106] is source 14: its Theorem III.36.1 (the low-degree symmetric threshold) and its Question 7 (the global quintic problem). “A specific question left open by the preceding report” is source 18’s Question 5 on $d = 1$, answered here in its leading coefficient (item II.38). On the heading “The direct Section 15 improvement”: the improvement is in the finite bound at moderate p ; at Lemma 15.5 the thresholds of Section II.84 gain nothing over source 18’s Corollary II.66.9 at leading order (note after Corollary II.84.5).

F.27.1 The direct Section 15 improvement

A d -arrangement consists of $2d$ parallel k -cubes. Its parameter space has dimension

$$D = (2d + 1)k + 2d - 1.$$

Gowers' Lemma 15.4 bounds the number of degenerate arrangements by $k3^{2d^2}p^{D-1}$. The preceding project analysis replaced its leading coefficient by $k + C_d$, where

$$C_d = \frac{1}{2} \left(T_{2d} - 2 \binom{2d}{d} + 1 \right), \quad T_{2d} = \sum_{a=0}^d \binom{2d}{2a} \binom{2a}{a}. \quad (\text{F.27.1})$$

It still bounded all nonparity certificates together by a very large multiple of p^{-2} .

The new estimate distinguishes their degrees. The degree-one contribution has coefficient $J_{k,d}$, given by a constant-term formula in Theorem II.83.2. For each fixed d it is a polynomial in k of degree $2d$. Certificates of higher degrees are assigned progressively smaller probabilities. The resulting finite bound is Theorem II.84.1. Its proof includes the necessary summation by parts; subtracting two unrelated cumulative upper bounds would not be a valid way to bound an individual degree stratum.

The case $d = 1$ has different geometry. We identify all its divisorial exceptions and prove

$$\delta_{k,1}(p) = \frac{k^2 + 3k}{p} + O_k(p^{-2}),$$

with an explicit finite remainder. This answers a specific question left open by the preceding report, including the exact formula $\delta_{1,1}(p) = 4/p - 3/p^2$.

F.27.2 Exact pointwise polynomial questions

Let Π prescribe which colors on a finite parameter set must agree. Given a polynomial coloring $Q : \mathbb{F}_q^n \rightarrow \mathbb{F}_q^r$, write $S_z(Q)$ for the directions h realizing that pattern from the base point z ; the zero direction is included. We prove an exact formula

$$\min_Q |S_z(Q)| = q^{\max\{n-B,0\}}. \quad (\text{F.27.2})$$

The integer B is computed by descending Gaussian elimination on a small matrix of parameter powers. This is an exact extremal answer at every *prescribed* center: the minimizing polynomial is allowed to depend on that center. It is generally a different question to minimize the sum over all centers using one polynomial.

For symmetric $\ell = 2m$ term progressions and degree caps D_a , put $s_a = \lceil D_a/2 \rceil$ and $t_a = \min(m, s_a)$. The answer uses

$$B_{\text{sym}} = \sum_a t_a (2s_a - t_a).$$

For monochromatic ℓ -term progressions it uses

$$B_{\text{mono}} = \sum_a \frac{r_a(2D_a - r_a + 1)}{2}, \quad r_a = \min(D_a, \ell - 1).$$

Both statements hold for arbitrarily large degree caps in characteristic $p > \ell$; a bound on the field size alone would be insufficient.

For symmetric four-term progressions, the scalar quintic and sextic pointwise budget is exactly eight. We exhibit the sharp example $N_3 + N_5$ at the origin. Thus a universal prescribed-center theorem cannot lower eight to seven. We also prove concrete obstructions to two natural ways of promoting such local examples to globally avoiding colorings.

F.27.3 What the claims mean

Status	Content
Established here	The stratified finite degeneracy bound; the balanced degree-one count and its growth; the $d = 1$ leading coefficient; the exact arbitrary-pattern pointwise minimum; the norm-extension obstructions and their counting forms.
Classical ingredients	Boolean interpolation, finite-field linear algebra, norm forms, and minimum support of polynomial functions, including its equality case.
Prior project inputs	The $d \geq 2$ leading coefficient $k + C_d$, the paired rank argument, and low-degree symmetric norm layers. Their required proofs are included.
Still unresolved	The actual second degeneracy coefficient in general; the global scalar quintic avoidance dimension; improved interval-coloring exponents from these ideas; propagation to a new global Szemerédi bound.

“Established here” describes mathematical proofs in this manuscript. It does not certify independent priority in the literature. The exact computations are reproducibility checks, not substitutes for those proofs. No new Lean kernel certification is claimed. The related norm-coloring work of Shi and Dong [33], and the classical results of Warning and Delsarte–Goethals–MacWilliams [40, 56, 96, 57], are explicitly credited.

F.28 Source 23. Results, provenance, and their place in the proof

Source 23, delivered Section 1.

Added in the fifth write, 6 October 2026. Delivered title page. “ProveIt / Combinatorics and Ramsey theory / Sharp Thresholds and Extremal Constants in Gowers’s Szemerédi Framework / Numerical semigroups, characteristic-five cube asymptotics, and common partitions for affine fibres / Research article prepared for integration into the ProveIt repository / 6 October 2026”; the PDF author field reads “Research note prepared for ProveIt”. Archive `gowers_sharp_thresholds.zip` (arrival 26908b3b6), manuscript `gowers_sharp_thresholds.tex` (2,320 lines, 33 pages), pin 65f4e039f (its `data/source_manifest.json`, shipped as `data/23-sharp-thresholds-source_manifest.json`). Printed below: the abstract, the closing lines of its title page and Section 1; Sections 2–6 are in Part I, Section 7 and Appendix A in Part III (a second route to source 21’s theorem, in its chapter, Section III.49), Section 8 in Part II (Section II.81), Sections 9–10 and Appendix B in Part V. Its macro `\snapshot` (its own pin) is renamed `\snapshottwentythree`; its explicit equation tags (F1), ... are kept.

Abstract of source 23 (as delivered). We establish precise local refinements and sharpness results for three parts of Gowers’s proof of Szemerédi’s theorem. First, simultaneous linear-phase partitions with cell lengths in a finite set $\mathcal{S}$ are controlled by the numerical-semigroup conductor

$C(\mathcal{S})$. A torus-packing upper bound and a mixed-radix lower construction give asymptotically matching constants when the permitted lengths are close. In particular, the eventual universal threshold for cell lengths $L, L+1$ is $L^{q+2}/\prod_i \varepsilon_i (1+o(1))$. This proves the necessity of the extra tiling cost left open in an earlier project report, even when different cells may use different steps. We also compute an exact exceptional-set profile for the obstructions and a sharp almost-partition threshold.

Second, for real centered functions on groups of exponent five, we prove the sharp cube-counting expansion

$$\|\delta + f\|_{U^3}^8 \leq \delta^8 + 6\sqrt{2}\delta^4 u^4 + \frac{2^{3/4}}{3}\delta^2 u^6 + O(u^8), \quad u = \|f\|_{U^3} \ll \delta,$$

with an absolute, dimension-independent error constant. An exact five-point model and indicator approximations attain the secondary coefficient; a rigidity theorem determines the necessary harmonic phase alignment. Third, a localization to width at least $m/3$ lets common slice refinements extend to full product boxes without further width loss. Paired interpolation and retained observations give an explicit conditional replacement for the geometric part of Lemma 16.10.

Every principal new statement is proved in the article. Exact finite certificates accompany the proofs. We distinguish these local results from the remaining global parameter comparison and from Lean formalization; no new general bound for Szemerédi’s theorem is claimed.

Source revision: 65f4e039fc4ad30a8f492e63abde2bcd0740d7c9.

Status: research draft with complete written proofs and reproducible finite checks.

Gowers’s 2001 proof develops several interacting quantitative mechanisms: uniformity norms control configuration counts; phase approximation turns algebraic structure into useful progressions; and repeated structural refinements make fibre information compatible across coordinates [1]. These mechanisms contain local questions whose answers can be sharpened independently of the full density-increment iteration. The present article concentrates on three such questions and proves explicit results for each.

The reference repository already contains extensive earlier refinements, and several changes to its transcription address the difference between the printed paper and statements suitable for formalization [107, 108, 109]. We treat those results as prior project work. In particular, the previously obtained upper exponent for prescribed linear-phase cell lengths is reproduced with attribution; the new contribution is a matching obstruction, a sharp leading constant, and the extension to arbitrary admissible length sets. Similarly, the leading centered three-cube coefficient was known in the project; here the order-five secondary coefficient is determined.

F.28.1 Three principal conclusions

Exact lengths have an unavoidable arithmetic cost. Let $T_{\mathcal{S}}(\varepsilon)$ be the eventual universal interval threshold for a partition into progressions of lengths in $\mathcal{S}$, on which q affine phases have arc diameters at most ε_i . If $L = \min \mathcal{S} \geq 2$, $H = \max \mathcal{S}$, $\gcd(\mathcal{S}) = 1$ and $C = C(\mathcal{S})$, then

$$C \prod_i \left(\left\lceil \frac{L-1}{\varepsilon_i} \right\rceil - 1 \right) \leq T_{\mathcal{S}}(\varepsilon) \leq C \left(\left\lceil \frac{(H-1)^q}{\prod_i \varepsilon_i} \right\rceil - 1 \right). \quad (\text{F.28.1})$$

For $\mathcal{S} = \{L, L+1\}$ this yields $T \sim L^{q+2}/\prod_i \varepsilon_i$, uniformly in the tolerances for fixed q as $L \rightarrow \infty$. This answers the prescribed-length sharpness question posed in [111, Question 11.1]. The lower example allows cell-dependent steps and persists for phases coming from a prime cyclic ambient group.

Order five selects a different secondary coefficient. For $G = \mathbb{F}_5^n$, $\mathbb{E}f = 0$, and $u = \|f\|_{U^3} \ll \delta$, the sharp first two nonconstant terms of the weighted cube supremum are

$$6\sqrt{2}\delta^4 u^4 \quad \text{and} \quad \frac{2^{3/4}}{3}\delta^2 u^6.$$

The proof retains the positive cost of the doubled dominant frequency and controls the incomplete cubes by the fourth-power Fourier tail. This turns the previous order estimate into a sharp coefficient with an $O(u^8)$ remainder. It also forces the dominant Fourier phase θ to satisfy $\cos(5\theta) \rightarrow 0$ in any family attaining the secondary term.

A common base refinement can keep its full width. For a proper prime-field box $T \times J$ of width at least m , a preliminary partition of T into short index boxes leaves width at least $m/3$. Any subsequent common-step base refinement of width $W > 1$ extends to a true partition of its product with J into boxes of width at least W . This uses full residue chains on the final axis. Combined with common slice covers and affine interpolation, it gives complete geometric and sampling conclusions under explicit inputs. The final Section 16 comparison of control functions remains a separate problem.

Table F.28.1: What is strengthened or settled in this article.

Connection	Previous project status	Result here
Lemma 5.9, linear phases	Conductor upper bound; necessity for prescribed lengths left open	Matching obstruction, leading constant one, and arbitrary admissible lengths
Exact versus partial partitions	Complete and flexible partitions studied separately	Exact packing formula for the obstruction; sharp small-loss threshold for one exact length
Lemma 3.10, three-cubes	Leading centered term and an order-five example; secondary optimum open	Secondary coefficient $2^{3/4}/3$, $O(u^8)$ error, and necessary phase rigidity
Lemma 16.10 geometry	Slice interpolation tools and partial transport already available	A common partition, exact width transport, and a conditional full-product covering theorem

F.28.2 Scope of the mathematical and computational claims

The article proves finite theorems and local asymptotics, including two answers to questions explicitly posed in the preceding project reports. We do not assert historical priority for all of the combined formulations. Dirichlet approximation, numerical-semigroup conductors, cube Cauchy–Schwarz, and two-point affine interpolation are established ingredients. The new arguments and the precise comparison with the earlier reports are identified where they occur.

These results do not constitute a new best general bound for $r_k(N)$, the largest size of a subset of $\{1, \dots, N\}$ with no nonconstant k -term arithmetic progression. For context, Leng, Sah, and Sawhney proved the stronger general estimate $r_k(N) \ll N \exp(-(\log \log N)^{c_k})$ for $k \geq 5$ in their 2024 paper [2]. Our local results are not compared against that bound as if they were an alternative global theorem.

The repository snapshot is 65f4e039fc4ad30a8f492e63abde2bcd0740d7c9, retrieved on 6 October 2026. Its status record explicitly separates completed formal components from unresolved statements and control-function comparisons [109]. We inspected the relevant source and proof interfaces but did not compile a new Lean development. The accompanying Python programs check finite arithmetic and polynomial identities; the proofs in the article establish the quantified mathematical claims.

[write] 6 October 2026; sources 21–27. *Identifications and credit within this report.* The “earlier project report” [111] is source 13: its Theorem 3.8 and Corollary 3.9 are Theorem I.26.8 and Corollary I.26.9, and its Question 11.1 is Research question V.54.1, answered below (Section I.38; item I.12). [112] is source 14 (Section III.35). The “research collection” [108] at this pin is the first write of this report. Section III.58 proves the theorem of source 21 independently (source 21’s Theorem F.26.1); in this report it is a credited second route inside source 21’s chapter. The sampling lemma of Section II.103 is measured against source 19’s theorem in the note after it.

F.28.3 Reading guide

Sections I.35–I.39 form a self-contained treatment of the phase-partition problem, including exact witnesses and algorithms. Section III.58 is independent of the partition argument and proves the characteristic-five cube result. Section II.103 supplies the geometric gluing theorem and its conditional Section 16 specialization. The final sections give a verification record, integration recommendations, and further research questions.

For reference, on a finite abelian group we use normalized averages and the Gowers norm convention

$$\|f\|_{U^s}^{2^s} = \mathbb{E}_{x, h_1, \dots, h_s} \prod_{v \in \{0,1\}^s} \mathcal{C}^{|v|} f(x + v \cdot h),$$

where $\mathcal{C}$ is complex conjugation. Fourier coefficients are $\hat{f}(\xi) = \mathbb{E}_x f(x) \overline{\xi(x)}$. Fourier sums have no normalizing factor. For real functions the conjugations in the physical-space cube product are immaterial. The cube Cauchy–Schwarz inequality used below is Gowers’s Lemma 3.8 [1]; in normalized form it bounds a mixed cube average by the product of the U^s norms of its 2^s vertex functions.

F.29 Source 24. Scope, headline results, and comparison with the source

Source 24, delivered Section 1.

Added in the fifth write, 6 October 2026. Delivered title block. “Sharper local structure in Gowers’s Szemerédi argument / Moment-sensitive dependent random choice, small difference sets, and vertical-fibre compression / Research draft prepared with ChatGPT / For the ProveIt research programme / 6 October 2026”; the PDF author field reads “Research draft prepared with ChatGPT”. Archive `gowers_section7_refinements.zip` (arrival e2bf630b0), manuscript `gowers_section7_refinements.tex` (1,711 lines, 22 pages), pin 2e0553b0c (its `PROVENANCE.md`, shipped as `24-section7-compression-PROVENANCE.md`). Printed below: the abstract, the verification-status statement after it and Section 1; Sections 2–9 form a chapter of Part II (Section II.72), Sections 10–13 and Appendices A–B are in Part V. Its macro `\code` (verbatim file names) is renamed `\codetwentyfour`; its ‘question’ environment (“Question”) is set in this report’s ‘Question’ style.

Abstract of source 24 (as delivered). We develop explicit, self-contained refinements of the structural reductions in Section 7 of Gowers’s *A new proof of Szemerédi’s theorem*. First, a cubic

supporting-line argument replaces the five-sample intersection lemma by a three-sample theorem: the retained fraction $\delta^5/\sqrt{2}$ becomes $7\delta^3/(3\sqrt{6})$, with the same 90% good-pair conclusion and the same intersection threshold. We determine the sharp constant in the corresponding scalar moment relaxation and identify exactly why its two-sample analogue fails. Second, retaining the actual edge mass of the popular-difference graph through a one-sample argument yields an elementary energy-to-structure bound with a linear retained fraction and relative doubling $2^{17}\kappa^{-6}$. Third, an injective vertical-fibre argument replaces the loss $8kC^{4k}$ in the Freiman-restriction lemma by kC^{2k+1} . The two structural improvements combine to replace the order-eight restriction bound $2^{-1882}\gamma^{1164}|B_0|$ by $2^{-295}\gamma^{103}|B_0|$. We also prove weighted and higher-arity selection statements, a finite-field-vector codomain extension, profile-dependent certificates, and an optimized threshold formula. The article separates these local improvements from stronger published Balog–Szemerédi–Gowers results and from global quantitative Szemerédi bounds. It includes complete proofs, a repository-specific formalization plan, exact finite regression tests, and further research questions. No new global record or literature-priority claim is made.

Verification status. The general results are proved mathematically below. The accompanying Python program performs exact finite checks; it is not a formal proof. No new Lean development was compiled for this article.

The question addressed here is not whether the entire 2001 proof should be replaced by contemporary higher-order Fourier analysis. It is more specific: which losses in its elementary structural machinery can be removed while retaining transparent statements, explicit constants, and a short dependency chain suitable for formalization? The answer developed below concerns Gowers’s Lemma 7.4, Proposition 7.3, Lemma 7.5, and Corollary 7.6 [1]. The corresponding statement and proof modules in the supplied repository were also inspected [113].

F.29.1 The compiled improvement

For a function $\phi : B_0 \rightarrow H$ between subsets of abelian groups, a *respected additive quadruple* is an ordered tuple $(a, b, c, d) \in B_0^4$ for which

$$a + b = c + d, \quad \phi(a) + \phi(b) = \phi(c) + \phi(d).$$

Repeated entries are allowed. A Freiman homomorphism of order k preserves every equality between two sums of k domain elements, again with repetitions allowed.

Theorem F.29.1 (Explicit order- k restriction). *Let G be an abelian group, p a prime, $B_0 \subset G$ a nonempty finite set of size m , and $\phi : B_0 \rightarrow \mathbb{F}_p$ a function. Suppose that at least γm^3 additive quadruples are respected, where $\gamma > 0$. For every integer $k \geq 1$, there is a subset $B \subseteq B_0$ on which ϕ is a Freiman homomorphism of order k and*

$$|B| \geq \frac{2^{-(34k+20)}}{k} \gamma^{12k+7} m. \tag{F.29.1}$$

In particular, for $k = 8$,

$$|B| \geq 2^{-295}\gamma^{103}|B_0|. \tag{F.29.2}$$

[write] **6 October 2026; sources 21–27.** Weaker than source 10 by 2^{10k+6} (checked at the write). Source 10’s Corollary II.21.2 gives, under the same hypotheses, $|B| \geq 2^{-(24k+14)}\gamma^{12k+7}m/k$; the ratio to the bound here is $2^{(34k+20)-(24k+14)} = 2^{10k+6}$, which is 2^{86} at $k = 8$ ($2^{-209}\gamma^{103}$ against $2^{-295}\gamma^{103}$). Both come from a Proposition 7.3-type step followed by a fibre extraction with $Q_k \leq C^{2k+1}$; source 10’s energy step

has the better constants. Like source 10’s, this bound implies the conclusion of the formal `corollary_7_6` (`Sections06_07.lean`, line 197) where the hypotheses agree, since $\gamma \leq 1$ gives $2^{-295}\gamma^{103} \geq 2^{-1882}\gamma^{1164}$; not formalized.

The hypotheses imply $\gamma \leq 1$: a graph with m points has additive energy at most m^3 . The theorem therefore improves the source’s lower bound throughout its meaningful parameter range. Its proof is completed in Section II.80; no asymptotic qualification such as “ p sufficiently large” is needed.

F.29.2 A ledger of the local changes

In the following table, $m = |A|$, $E(A) \geq \kappa m^3$, and $|\Gamma - \Gamma| \leq C|\Gamma|$ for the graph in the restriction lemma. The energy parameter satisfies $0 < \kappa \leq 1$, and $C \geq 1$ for a nonempty graph. The source’s Proposition 7.3 is stated for subsets of $\mathbb{Z}^n$; the proof here works in every abelian group.

Target	Source guarantee	Guarantee proved here
Lemma 7.4: core size	$\delta^5 n / \sqrt{2}$	$7\delta^3 n / (3\sqrt{6})$
Proposition 7.3: retained size	$2^{-20}\kappa^{12}m$	$2^{-3}\kappa m$
Proposition 7.3: difference set	$2^{38}\kappa^{-24}m$	$2^{15}\kappa^{-5}m$
Lemma 7.5: retained size	$ B /(8kC^{4k})$	$ B /(kC^{2k+1})$
Corollary 7.6: order-eight size	$2^{-1882}\gamma^{1164}m$	$2^{-295}\gamma^{103}m$

Table F.29.1: Comparison with the specified Section 7 statements, not a table of current world records. The good-pair fraction and threshold in the first row are unchanged.

Two different selection regimes appear. The three-sample theorem strictly strengthens Lemma 7.4 at its original parameters. The strongest explicit energy bound in this article instead uses *one sample* with a lower codegree threshold. These are parallel applications of the same exact moment identity; the proof of the exponent 103 does not silently rely on substituting the three-sample bound into every later line of Gowers’s proof.

[write] **6 October 2026; sources 21–27.** *Within this report most of the table repeats earlier sources* (checked at the write). This source’s pin (2e0553b0c, 13:44) predates the placement of sources 03 and 10, so the repetition is independent; but in this report: the first row is source 03’s Theorem II.5.2 (the same certificate, tangent at $7u/6$); the second and third rows, $(2^{-3}\kappa, 2^{15}\kappa^{-5})$, lie between source 03’s $(c/8, 2^{19}c^{-9})$ (Theorem II.6.3) and source 10’s $(c/4, 2^{10}c^{-5})$ (Theorem II.20.1); the fourth row’s $|B|/(kC^{2k+1})$ is source 10’s Theorem II.21.1, whose sharper form $|B|/(k(h_k - 1) + 1)$ is never worse than Theorem II.79.3’s $|B|/(kQ_k)$; and the last row is weaker than source 10’s $2^{-209}\gamma^{103}$ (Corollary II.21.2; note after Theorem F.29.1). “The source’s Proposition 7.3 is stated for subsets of $\mathbb{Z}^n$; the proof here works in every abelian group”: the group form is already formal at HEAD 2dfe46f44, with the paper’s constants, as `proposition_7_3_group_holds` (`Proofs07BalogSzemerédi.lean`, line 854; it existed at this source’s pin), beside `proposition_7_3_holds` (line 996) for $\mathbb{Z}^n$; only the constants here are new, and source 10’s are better. New in this report: Theorem II.79.5, the selector of Section II.73, and the threshold optimizations of Sections II.77 and II.80.

F.29.3 What is classical, what is derived, and what is not claimed

Dependent random choice, popular-difference graphs, path counting, and Plünnecke–Ruzsa growth inequalities are established tools [1, 11, 36]. Our contribution here is an explicit reorganization and sharpening of these particular interfaces, including the scalar certificates and the vertical-fibre

bookkeeping. The proofs establish the stated inequalities independently of any assertion that their individual formulations have not appeared elsewhere.

Stronger published Balog–Szemerédi–Gowers theorems already exist. For example, Reiher and Schoen obtain a subset of size $(1 - \varepsilon)\kappa^{1/2}|A|$ with relative difference-set bound $C_\varepsilon\kappa^{-4}$ [12]. We explicitly import that theorem in Section II.80.3 to derive a further restriction estimate, with its dependence on the imported theorem visible. Thus the elementary energy bound here is not advertised as a new extremal record.

Nor is a local improvement automatically a new quantitative Szemerédi bound. As a published modern benchmark, Leng, Sah, and Sawhney prove, for every fixed $\ell \geq 5$, a bound of the form

$$r_\ell(N) \ll N \exp(-(\log \log N)^{c_\ell}) \quad (c_\ell > 0)$$

by substantially different machinery [2]. We do not re-run Gowers’s Sections 8–18 with the constants below, and we make no claim about the best currently available global estimate. The sharpness result proved here concerns a precisely defined scalar relaxation, not an unrestricted graph or additive-combinatorial extremal problem.

F.30 Source 25. Scope, main theorem, and relation to the source

Source 25, delivered Section 1.

Added in the fifth write, 6 October 2026. Delivered title block. “The Codimension-One Degeneracies of Gowers Arrangements / An asymptotically sharp replacement for Lemma 15.4 and explicit random-selection thresholds / Research manuscript prepared for Vladimir Reshetnikov / Mathematical development and computational checks with ChatGPT / 6 October 2026”; *the PDF author field reads* “Research prepared with ChatGPT for Vladimir Reshetnikov”. Archive `gowers_arrangement_degeneracy.zip` (arrival `d7c4aa399`), manuscript `arrangement_degeneracy.tex` (1,207 lines, 17 pages), pin `568d1802e` (its `SOURCE_AUDIT.md`, shipped as `25-codimension-one-SOURCE_AUDIT.md`). Printed below: the abstract and Section 1; Sections 2–8 are in Part II (Section II.81), Sections 9–11 in Part V. Its macro `\card` (an absolute value, not this report’s operator) is renamed `\cardtwentyfive`; its ‘question’ environment (“Question”, numbered on its own counter there) is set in this report’s ‘Question’ style, numbered within its section (its Questions 1–8 keep their indices).

Abstract of source 25 (as delivered). We study the full ternary-moment degeneracy condition in Section 15 of Gowers’s proof of Szemerédi’s theorem. For a d -arrangement of k -cubes, put $D = (2d + 1)k + 2d - 1$ and $n = 2d2^k$. Over a finite field of cardinality q and characteristic greater than $2d$, with $k \geq 1$ and $d \geq 2$, we prove that the number of degenerate labelled arrangements is

$$(k + H_d)q^{D-1} + O_{k,d}(q^{D-2}), \quad H_d = \frac{T_{2d} - 2\binom{2d}{d} + 1}{2},$$

where T_m is the central trinomial coefficient. All constants in the error term are explicit. In particular, $H_8 = 2\,585\,444$. The coefficient $k + H_d$ is exact: the codimension-one components are precisely the zero sidelength hyperplanes and a projective family of balanced ternary cross-section hyperplanes. Outside those hyperplanes, every remaining witness has codimension at least two. The key step is a minimal-Walsh-layer linearization, followed by use of the moment equations involving the last coordinate that the original counting argument does not need to retain. A sharp single-witness probability, a genuine second-order obstruction, and a small-characteristic counterexample accompany the classification. We also prove a finite-field, arbitrary- d random-selection theorem with exact Riesz-product parameters and a better explicit ambient-size requirement. The package includes reproducible exact finite checks and a ProveIt integration plan. These are local refinements, not a claimed improvement of the final quantitative Szemerédi theorem; no proof-assistant verification or literature-priority claim is made.

Gowers's Lemma 15.4 bounds exceptional arrangements before the random selection in Lemma 15.5. In the original notation, its estimate is

$$|\mathcal{D}_{k,d}(\mathbb{F}_N)| \leq k 3^{2d2^k} N^{(2d+1)k+2d-2}. \quad (\text{F.30.1})$$

The proof chooses one nonconstant multilinear moment for each ternary coefficient pattern, uses a polynomial fibre bound, and then takes a union bound [1, pp. 562–565]. The resulting constant treats all witnesses as though they defined hypersurfaces. This is the point we sharpen.

The full definition of degeneracy uses *every* coordinate monomial, including monomials that contain the last, cross-section coordinate. Gowers remarks after Lemma 15.4 that his argument could use fewer such monomials. For the present purpose, retaining them is decisive: they usually give a second independent equation.

The ProveIt interface inspected here is `Sections14_15.lean` at commit `568d1802e99404c0710be33cceb401d25211f445`. It defines degeneracy by a ternary coefficient function which is not a *modular* multiple of the distinguished parity coefficient, and requires all arrangement moments to vanish [114]. Our definition agrees with that interface, after an invertible change from Boolean to centred sign coordinates. The existing research README was also inspected to separate this project from its density-transfer, phase-partition, energy, cube-counting and floor-pattern developments [115].

F.30.1 Parameters

Throughout the main results,

$$k \geq 1, \quad d \geq 2, \quad q = |\mathbb{F}_q|, \quad p = \text{char } \mathbb{F}_q > 2d, \quad n = 2d2^k, \quad D = (2d+1)k + 2d - 1. \quad (\text{F.30.2})$$

Define

$$T_m = [z^0](z^{-1} + 1 + z)^m = \sum_{j=0}^{\lfloor m/2 \rfloor} \binom{m}{j} \binom{m-j}{j}, \quad (\text{F.30.3})$$

$$H_d = \frac{T_{2d} - 2\binom{2d}{d} + 1}{2}, \quad L_{k,d} = k + H_d, \quad U_{k,d} = \frac{3^{2d2^k} - 3^{2d}}{2}. \quad (\text{F.30.4})$$

The arrangement count is a count of labelled parameter tuples, including repeated vertices when these occur. There are exactly q^D such tuples satisfying the cross-section relation.

Theorem F.30.1 (Explicit sharp first-order count). *Under (F.30.2), the number $Z_{k,d}(q)$ of degenerate arrangements satisfies*

$$L_{k,d}q^{D-1} - \binom{L_{k,d}}{2}q^{D-2} \leq Z_{k,d}(q) \leq L_{k,d}q^{D-1} + U_{k,d}q^{D-2}. \quad (\text{F.30.5})$$

Consequently, for fixed k, d ,

$$\lim_{\substack{q \rightarrow \infty \\ \text{char } \mathbb{F}_q > 2d}} \frac{Z_{k,d}(q)}{q^{D-1}} = k + H_d. \quad (\text{F.30.6})$$

Over an algebraically closed field of characteristic zero or greater than $2d$, the codimension-one irreducible components of the same degeneracy locus are exactly $k + H_d$ explicitly described hyperplanes.

[write] 6 October 2026; sources 21–27. *Against the formal statement* (checked at the write). lemma_15_4 (Sections14_15.lean, line 402) bounds the count by $k3^n q^{D-1}$, $n = 2d2^k$, for prime q . Here $L_{k,d} + U_{k,d}/q \leq k3^n$ for $d \geq 2$, $q \geq 5$: $H_d \leq (3^{2d} + 1)/2 \leq (3^{n/2} + 1)/2$ because $n \geq 4d$, and $U_{k,d}/q < 3^n/10$, so $L_{k,d} + U_{k,d}/q < k + (3^{n/2} + 1)/2 + 3^n/10 \leq k3^n$ (already for $k = 1$, since $0.9 \cdot 3^n - 1 \geq (3^{n/2} + 1)/2$ for $n \geq 8$). So the theorem implies the formal bound for $d \geq 2$ and primes $q > 2d$, as source 18’s does; not formalized. Its predicate is the formal one (note in Section V.102).

The proof occupies Sections II.87–II.90. The leading constant cannot be reduced. For the 8-arrangements actually used by Gowers, (F.30.5) gives

$$Z_{k,8}(q) \leq (k + 2\,585\,444)q^{17k+14} + \frac{3^{2^{k+4}} - 3^{16}}{2} q^{17k+13}. \quad (\text{F.30.7})$$

Thus the enormous ternary-pattern count has moved from the leading q^{-1} term to a q^{-2} error term. This does not eliminate the large error constant, and it is not a dimension-uniform asymptotic.

F.30.2 What is, and is not, established

The main contributions developed here are the minimal-layer rank bound, the separation of parity and non-parity degeneracies, the exact projective hyperplane count, and their combination into (F.30.5). Section II.92 gives a self-contained quantitative application to the selection step. The hypotheses $d \geq 2$ and $p > 2d$ are substantive; we show what fails without them.

The proofs are conventional mathematical proofs, supplemented by finite exact computations. They have not been independently refereed or formalized in Lean. A targeted source search does not establish priority for the classification in the broader literature. In particular, the standard finite-field viewpoint on hyperplane arrangements is classical [98]; it is not presented here as a new method. Nor do we claim to have propagated these local savings through Sections 16–18 to obtain a new bound for progression-free sets.

[write] 6 October 2026; sources 21–27. *Credit and novelty within this report.* The leading coefficient $k + H_d$ with $H_d = (T_{2d} - 2\binom{2d}{d} + 1)/2$ is source 18’s Theorem II.66.8, with $H_d =$ source 18’s C_d , for $d \geq 2$ and primes $p > 2d$; the sharp single-witness bound (Theorem II.88.1) is source 18’s Proposition II.66.3. Source 18’s archive arrived at 14:45 and was not opened before this source’s pin (14:52), so the derivation here is independent; in this report source 18 is the first print and this is a second route. The “existing research README” [115] inspected here was that of the first write (sources 01–05). New in this report: prime-power fields, the classification of the codimension-one components over an algebraically closed field, the residual lower bound of Theorem II.90.1, the small-characteristic counterexample, and the selection theorem for every $d \geq 2$ (Theorem II.92.1). Source 27 re-derives the leading coefficient a third time and determines the second coefficient for $d = 2$.

F.31 Source 26. The question and the scope

Source 26, delivered Section 1.

Added in the fifth write, 6 October 2026. Delivered title block. “Threshold free affine density increments on cyclic groups / Report275 / Research report prepared for Vladimir / 6 October 2026”; the PDF author field reads “Research report prepared for Vladimir”. Archive Report275_Threshold_Free_Affine_Density_Increments.zip (arrival 7e827c6f3), manuscript article.tex (494 lines, 12 pages), pin c94998a5d (its SOURCES.md, shipped as 26-threshold-free-affine-SOURCES.md). Printed below: the abstract and Section 1; Sections 2–4 and Appendix A form a chapter of Part I (Section I.40), Sections 5–6 are in Part V.

Abstract of source 26 (as delivered). A family of local affine correlations on disjoint modular progressions covering a set yields a density increment with no scale threshold. We keep the exact balanced mass of the cover. If its shortest progression has length r , the set has density δ , the normalized balanced mass is μ , and the summed correlation is at least αN , then a proper progression has density at least $\delta + \alpha/3$ and length at least

$$\min \left\{ \left(\frac{r(1-\delta)}{128\mu} \right)^{1/3}, (N/8)^{1/3} \right\}.$$

A complementary, parent-contained progression has the same density increment and length at least $\sqrt{\alpha r/(24\pi\mu)}$, over every cyclic modulus. For every cyclic modulus and M pairwise two-comparable parents, the resulting guarantees are $(N/M)^{1/3}/8$ and $\sqrt{\alpha(N/M)/(96\pi(1-\delta))}$. This strengthens the full degree-one specialization of a specified repository version of Corollary 5.8. The proofs combine a subgroup-coset moment bound, removal of nonproper directions, and phase-flat residue partitions. A separate appendix records a common-step integer-box almost-cover, with its exact scope. The companion checks bounded identities and regressions; no formal verification, global Ramsey bound, or literature-priority claim is made.

The local density-increment step is a useful place to separate an actual quantitative improvement from a change of hypotheses. The target here is the degree-one part of the repository predicate `corollary_5_8` at the source pin specified in Section V.105. Its displayed premises do not include a lower-scale assumption. A phase-partition argument alone can demand a minimum parent length that is not supplied by those premises. We avoid this issue by using a second-moment increment when the requested correlation gain is small, and phase refinement when it is large.

The principal results below are self-contained written theorems over cyclic groups. Their strongest form uses the actual balanced mass on the cover and does not require comparable parent lengths. Comparability enters only to express the answer through N/M . All progressions in the main body are modular and injectively parameterized. Integer representatives are treated separately in Appendix I.44.

The cube-root scale, Dirichlet phase partitioning, and progression moments have classical precedents. Section V.105 gives a narrow source comparison. The claims concern the explicit local interfaces proved here; they do not close the general-degree catalogue statement, change the repository’s formal status, or improve a final Szemerédi or Ramsey bound.

[write] **6 October 2026; sources 21–27.** *The degree-one formal statement* (checked at HEAD 2dfe46f44; `Section05.lean`, `Definitions.lean` and `Proofs05Corollary58.lean` have the same blobs as at this source’s pin). `corollary_5_8` (`Section05.lean`, line 242) is open; at $k = 1$ it asks for a proper progression Q with $(N/M)^{1/16}/8 \leq |Q|$ and $(\delta + \alpha/16)|Q| \leq |A \cap Q|$. *Proved at the write:* that instance holds trivially whenever $N/M \leq 2^{48}$. The correlation hypothesis forces $A \neq \emptyset$ and $0 < \delta < 1$, and $\alpha N \leq \sum_{x \in U} |\mathbf{1}_A(x) - \delta| \leq 2\delta(1-\delta)N$ (the bound of (I.41.4)); a singleton $Q = \{a\}$ with $a \in A$ is a proper modular progression (`ModAP.IsProper` asks only that the carrier have the stated length), its density 1 is at least $\delta + \alpha/16$, and $(N/M)^{1/16}/8 \leq 1$ exactly when $N/M \leq 8^{16} = 2^{48}$. Source 26’s bound $(N/M)^{1/3}/8$ exceeds 1 only when $N/M > 512$; for the formal instance its content is the range $N/M > 2^{48}$, where it implies that instance (note after Corollary I.41.3). The ledger does not change, because `corollary_5_8` quantifies over every degree; the repository proves only the scale-qualified repair `corollary_5_8_with_scale_holds` (`Proofs05FullPartition.lean`, line 94). A formalization candidate; not formalized.

F.32 Source 27. The target and the main conclusions

Source 27, delivered Section 1.

Added in the fifth write, 6 October 2026. Delivered title page. “Combinatorics / Ramsey theory / Sharp Degeneracy Bounds for Gowers Arrangements / Hyperplane obstructions, second-order structure, and an exact finite-field count / An AI-assisted research draft prepared for Vladimir Reshetnikov and the ProveIt repository / 6 October 2026”; the PDF author field reads “AI-assisted research draft prepared for Vladimir Reshetnikov”. Archive `gowers_arrangement_degeneracy(2).zip` (arrival 6078f6e16), manuscript `article.tex` (1,824 lines, 28 pages); no commit pin: its `SOURCE_AUDIT.md` (shipped as `27-degeneracy-second-order-SOURCE_AUDIT.md`) and its bibliography record blob identifiers only. Printed below: the abstract, the status statement of its title page and Section 1; Sections 2–11 are in Part II (Section II.81), Sections 12–14 and Appendix A in Part V. Its `cleveref` references are written out as its own `\crefname` settings print them (capitalized names, equations as a number in parentheses); its ‘question’ environment (“Research question”, on the shared counter) is set in this report’s ‘Research question’ style.

Abstract of source 27 (as delivered). Gowers’s Lemma 15.4 bounds the proportion of degenerate labelled d -arrangements by $k3^{2d^k}/q$. We isolate the actual hyperplane obstructions and prove a sharp first-order asymptotic over every finite-field characteristic. In sufficiently large characteristic its coefficient is $k + L_d$, where

$$L_d = \frac{T_{2d} - 2\binom{2d}{d} + 1}{2}$$

and T_r is the central trinomial coefficient. Outside these hyperplanes, we classify $S_{k,d} = k(2^{2d} - 1) + k(k - 1)$ coefficient directions contributing at order q^{-2} , and bound all remaining directions by a cubic remainder. This replaces a single exponential union bound by separate linear, square-root, and cube-root field-size requirements.

For $d = 2$ the first two asymptotic coefficients are determined explicitly. For $k = 1, d = 2$, a finite rational incidence certificate gives the exact proportion, in every characteristic at least 11,

$$\frac{5}{q} + \frac{8}{q^2} + \frac{40}{q^3} - \frac{367}{q^4} + \frac{563}{q^5} - \frac{248}{q^6}.$$

The general theorems have conventional proofs. The last specialization also uses a completely specified finite enumeration, reproduced by the accompanying code and checked independently through subset sums of moment columns.

Status and scope. Unrefereed research draft; no proof-assistant verification or independent mathematical review is claimed. Priority for the refinements has not been established. These are local strengthenings of an exceptional-arrangement lemma, not a claimed new global bound for Szemerédi’s theorem. The all-squarefree-moments definition of degeneracy is essential.

Gowers’s proof of Szemerédi’s theorem uses higher-dimensional arrangements and a random selection step that favors arrangements respected by a partial map. Lemma 15.4 controls the exceptional arrangements on which the expected selection behavior can fail [1, §15, pp. 562–565]. This article studies that exceptional set itself. The aim is to replace a coarse count of possible coefficient arrays by a structural classification of their zero sets.

The relevant conventions are fixed in Section II.93. In particular, arrangements are *labelled parameter tuples*, not unlabelled geometric sets; repeated vertices are allowed. Write

$$n = 2d, \quad m = n2^k, \quad D = (n + 1)k + n - 1. \quad (\text{F.32.1})$$

There are exactly q^D arrangements in the full ambient space. Let $\Delta_{k,d}(q)$ be the fraction that are degenerate. Unless explicitly stated otherwise, $k \geq 1$ and $d \geq 2$.

The printed bound in Lemma 15.4 is

$$\Delta_{k,d}(q) \leq \frac{k3^m}{q} \quad (\text{F.32.2})$$

in its prime-field setting [1, p. 563]. It is obtained by charging roughly k/q to each ternary coefficient array. Most arrays do not define a hypersurface of bad arrangements. That observation is the source of the improvement.

F.32.1 Clean large-characteristic statements

Define the central trinomial coefficient and three explicit constants by

$$\mathsf{T}_r = [z^0](z^{-1} + 1 + z)^r = \sum_{j=0}^{\lfloor r/2 \rfloor} \binom{r}{2j} \binom{2j}{j}, \quad (\text{F.32.3})$$

$$L_d = \frac{\mathsf{T}_n - 2\binom{n}{n/2} + 1}{2}, \quad (\text{F.32.4})$$

$$R_{k,d} = \frac{\mathsf{T}_m - 2\binom{m}{m/2} - 3^n + 2^{n+1}}{2}, \quad (\text{F.32.5})$$

$$S_{k,d} = k(2^n - 1) + k(k - 1). \quad (\text{F.32.6})$$

The integrality and the meaning of these quantities will be proved, rather than assumed. L_d counts cross-section hyperplanes; $R_{k,d}$ counts non-parity coefficient directions; $S_{k,d}$ counts a distinguished second-order subfamily.

Theorem F.32.1 (Sharp first order and a cubic remainder). *Let $\mathbb{F}_q$ have characteristic $p > m = 2d2^k$. Put $A = k + L_d$. Then*

$$\frac{A}{q} - \frac{\binom{A}{2}}{q^2} \leq \Delta_{k,d}(q) \leq \frac{A}{q} + \frac{S_{k,d}}{q^2} + \frac{R_{k,d} - S_{k,d}}{q^3}. \quad (\text{F.32.7})$$

Consequently,

$$\Delta_{k,d}(q) = \frac{k + L_d}{q} + O_{k,d}(q^{-2}). \quad (\text{F.32.8})$$

The coefficient $k + L_d$ is sharp.

[write] **6 October 2026; sources 21–27.** *Against the formal statement* (checked at the write). lemma_15_4 (Sections14_15.lean, line 402) gives $\Delta_{k,d}(N) \leq k3^m/N$ for primes N , $m = 2d2^k$. For $p > m$ the upper bound here is at most $(k + L_d)/q + S_{k,d}/q^2 + R_{k,d}/q^3$ with $L_d \leq (3^{m/2} + 1)/2$, $S_{k,d} < k(2^{m/2} + k)$ and $R_{k,d} < (3^m + 2^{m/2+1})/2$; since $q > m \geq 8$ these give less than $k3^m/q$. So the theorem implies the formal bound for $d \geq 2$ and primes $p > m$, as source 18's does for $p > 2d$; not formalized.

The displayed lower bound may be negative for small q ; it can always be replaced by its maximum with 0. Likewise any upper bound can be capped at 1. The theorem is useful as a stratified estimate, not merely as a fixed-parameter asymptotic. The q^{-1} coefficient no longer grows exponentially in 2^k .

Corollary F.32.2 (An explicit exceptional-set budget). *Under the hypotheses of Theorem F.32.1, let $0 < \varepsilon \leq 1$. It is sufficient for $\Delta_{k,d}(q) \leq \varepsilon$ that*

$$q \geq \max \left\{ \frac{3(k + L_d)}{\varepsilon}, \sqrt{\frac{3S_{k,d}}{\varepsilon}}, \sqrt[3]{\frac{3(R_{k,d} - S_{k,d})}{\varepsilon}} \right\}. \quad (\text{F.32.9})$$

The characteristic hypothesis is separate from this size requirement.

Proof. Each of the three nonnegative summands in the upper bound in (F.32.7) is at most $\varepsilon/3$. $\square$

Theorem F.32.3 (Two pairs of cross sections). *For $d = 2$, $k \geq 1$, and characteristic $p > 4 \cdot 2^k$,*

$$\Delta_{k,2}(q) = \frac{k+4}{q} + \frac{k^2+21k-6}{2q^2} + O_k(q^{-3}). \quad (\text{F.32.10})$$

For $k = 1, d = 2$ the following exact formula holds over every finite field of characteristic at least 11:

$$\Delta_{1,2}(q) = \frac{5}{q} + \frac{8}{q^2} + \frac{40}{q^3} - \frac{367}{q^4} + \frac{563}{q^5} - \frac{248}{q^6}. \quad (\text{F.32.11})$$

The exact formula uses the finite incidence certificate in Section II.100; the asymptotic formula does not.

[write] **6 October 2026; sources 21–27.** *Answers source 18's first question for $d = 2$ (item II.38; also source 22's first and source 25's first question). Checked at the write: Corollary II.99.2 with $L_2 = 4$, $S_{k,2} = 15k + k(k-1) = k^2 + 14k$ and $b_{2,2,p} = 3$ (Proposition II.98.1) gives $S_{k,2} - \binom{k}{2} - 4k - 3 = (k^2 + 21k - 6)/2$, that is 8, 20, 33, 47 for $k = 1, \dots, 4$; at $k = 1$ this is the $8/q^2$ of the exact formula. The intake confirmed the exact formula by an independent exhaustive count at $q = 11, 13, 17, 19, 23$ (a different normalization and a meet-in-the-middle count of all ternary relations), and found it fails at $q = 5, 7$, as the source says. A Monte Carlo at $k = 2$ ($q = 53$ and 101) found the core part near $S_{2,2}/q^2 = 32/q^2$ and the rest shrinking like q^{-3} ; that is evidence for the case $k = 2$, not a proof. The asymptotic needs characteristic $p > 4 \cdot 2^k$ as stated; its proof needs only $p > 2k + 2$ (the remark in Section II.97).*

F.32.2 What is proved, and what is not

The principal contributions are the classification of the first-order obstructions, the exact second-order coefficient-direction count, the cubic-remainder bound, and the exact eight-vertex specialization. The all-characteristic extension in Section II.96 is not a formal consequence of the original odd-prime Walsh argument: it uses Boolean coefficient extraction, which is valid in characteristic two as well.

The argument does not change the definition of an arrangement, claim a new inverse theorem, or rerun the whole proof of Szemerédi's theorem. In particular, reducing one exceptional-set threshold does not by itself prove a better final density bound. The proposed questions in Section V.108 separate further local problems from that global bookkeeping task.

The repository source inspected for compatibility is `Combinatorics/Ramsey/Lean/GowersSzemerédi/Sections14_15.lean` [120]. It defines the moments, modular-multiple exclusion, and Lemma 15.4 proposition used here. A proposition-valued statement in that file is not evidence that the new theorems in this article have been formalized. A precise integration map appears in Section V.107.

[write] **6 October 2026; sources 21–27.** *Credit and novelty within this report.* The first-order coefficient $k + L_d$, with $L_d =$ source 18's C_d , and the per-direction bound of Proposition II.94.3 are source 18's Theorem II.66.8 and Proposition II.66.3, printed first; this source has no pin and cites neither source 18 nor source 22, whose texts were not in the repository, so its derivation is independent, and a third route beside source 25's. Its sentence in Section V.107 that the research directory “already contains numerous local refinements” and that this article “focuses on the degeneracy count rather than reproducing” them overlooks sources 18 and 22. New in this report: the first order in every characteristic (Theorem II.96.3), the core-direction classification with a cubic remainder (Theorem II.97.8), the second coefficient for $d = 2$

(Theorem F.32.3), which answers source 18’s first question for $d = 2$, the reduction of general d to the statistic $b_{2,d,p}$ (Corollary II.99.2), and the exact eight-vertex count. Its endpoint $k = d = 1$ (Proposition II.102.1) is source 22’s Theorem II.85.1 at $k = 1$.

F.33 Source 28. Results and the scope of the improvement

Source 28, delivered Section 1.

Added in the sixth write, 6 October 2026. Delivered title page. “ProveIt research report / Logarithmic Selection, Local Freiman Models, and Higher Uniformity / Refinements of Gowers’s Szemerédi framework / Prepared for Vladimir Reshetnikov / Mathematical development and computational checks with ChatGPT / 6 October 2026”; *the PDF author field reads “Research report prepared with ChatGPT for Vladimir Reshetnikov”*. Archive `gowers_logarithmic_selection.zip` (arrival 624d7a42a), manuscript `gowers_logarithmic_selection.tex` (2,594 lines, 36 pages), pin `e3125e698` (its `data/source_provenance.json`, shipped as `data/28-log-selection-source_provenance.json`). Printed below: the abstract, the boxed principal conclusions and closing lines of its title page, and Section 1; Sections 2–7 are in Part II (Section II.121), Section 8 in Part III (Section III.66), Sections 9–12 in Part V. Its macro `\T` (the circle $\mathbb{R}/\mathbb{Z}$) is set as this report’s `\Tcirc` and its `\pin` as `\pintwentyeight`; its `\cleveref` references are written out as `\cleveref` prints them there (lowercase names, “equation (4.3)”); its ‘question’ environment (“Research question”, on its own counter) is set in this report’s ‘Research question’ style; its Research questions 1–15 keep their indices.

Abstract of source 28 (as delivered). We refine two local questions arising from W. T. Gowers’s proof of Szemerédi’s theorem and the accompanying ProveIt research collection. The main result replaces the component selection in Section 10 by random windows in the logarithms of the fibre size and the anchored additive-energy weight. Boundary losses are measured with those same weights. This retains the correlation between useful mass and exceptional mass throughout the selection argument. We then normalize the anchored agreement directly, choose modal directional differences, and use coupled paths to obtain an exact Freiman model.

For a domain of density α and an order-eight relation defect at most 2^{-21} , the resulting theorem retains at least $\alpha(\log 2)^2/[20 \log(32/\alpha) \log(32/\alpha^2)]$ of the original points. Its spectrum has size at most $2^{34}(1 + \alpha)/[(\log 2)^2 \alpha^3]$, and its Bohr radius is at least $(\log 2)^2 \alpha^5/[3 \cdot 2^{69} \pi(1 + \alpha)^2]$. The source group may be any finite abelian group and the target any abelian group. All frequencies come from the original domain. The local difference map preserves every order-seven relation. Earlier repository work supplies the conceptual precedents for coupled paths, dense Freiman extension, and Bohr regularization; these ingredients are reproved here, and their provenance is separated from the new selection and parameter chain.

Independently, we examine the open higher-dimensional analogue of the repository’s sharp odd-order U^3/U^2 inequality. An exact two-harmonic polynomial proves that adding a small negative third harmonic decreases the U^4/U^2 ratio of a cosine. A rational parameter gives $c_4^{\text{odd}} > 0.55872939$, improving the single-cosine lower benchmark 0.51813359. We also prove $c_d^{\text{odd}} \leq \sqrt{160/323} = 0.70381534 \dots$ for every $d \geq 4$, and extend the failure of single-cosine local extremality to every $d \geq 4$. The sharp constants remain open. We provide reproducible finite checks, integration interfaces, and further research questions.

Principal conclusions. An order-eight approximate homomorphism on a finite abelian domain admits an exact local model determined by the Fourier coefficients of the original domain. The guaranteed retained fraction is

$$\frac{\alpha(\log 2)^2}{20 \log(32/\alpha) \log(32/\alpha^2)},$$

the number of defining frequencies is $O(\alpha^{-3})$, and the guaranteed Bohr radius is $\Omega(\alpha^5)$. The allowed proportion of failed order-eight relations is 2^{-21} .

Exact norm calculations give a new interval for the odd-order U^4/U^2 comparison constant and prove that a single cosine is not locally extremal at any order $d \geq 4$ when the character order is sufficiently large.

Repository comparison pinned to [e3125e6982abe2d44bd7cb9369063d4e69d41a2d](#). Complete proofs, exact finite certificates, and a proposed research agenda. Research draft; no new Lean verification or global Szemerédi bound is asserted.

The task is to strengthen actual steps of Gowers’s argument while retaining the information needed by later steps. Section 10 is a particularly useful target: the set of frequencies in its conclusion must be extracted from the original domain, not from an arbitrarily selected structured subset. Gowers’s Theorem 10.13 is designed to preserve exactly that information [1]. The theorem below has the same purpose, with different explicit controls.

F.33.1 The main local-model theorem

Let G be a finite abelian group, of cardinality N . A *finite domain over G* consists of a nonempty finite set X and a map $r : X \rightarrow G$. Put $X_s = r^{-1}(s)$. An order- k approximate homomorphism $\phi : X \rightarrow H$, where H is abelian, respects a prescribed proportion of the ordered $2k$ -tuples whose two base sums agree. Exact definitions appear in section II.122. Fourier coefficients in the next theorem are unnormalized.

Theorem F.33.1 (A spectrum-defined exact local model). *Let G be finite abelian, let H be abelian, and suppose*

$$0 < \alpha \leq 1, \quad M > 0, \quad |X_s| \leq M, \quad |X| = \alpha MN.$$

Suppose $\phi : X \rightarrow H$ respects at least a fraction $1 - 2^{-21}$ of the order-eight relations of the domain. Define

$$p(\alpha) = \frac{(\log 2)^2}{\log(32/\alpha) \log(32/\alpha^2)}, \tag{F.33.1}$$

$$\lambda(\alpha) = 2^{-17} \alpha^2 \sqrt{\frac{\log 2}{1 + \alpha}}, \tag{F.33.2}$$

$$g(s) = |X_s|, \quad \Gamma = \{\chi \in \widehat{G} : |\widehat{g}(\chi)| \geq \lambda(\alpha) MN\}, \tag{F.33.3}$$

$$R_\alpha = \frac{2^{34}(1 + \alpha)}{(\log 2)\alpha^3}, \tag{F.33.4}$$

$$\zeta_\alpha = \frac{(\log 2)^2 \alpha^5}{3 \cdot 2^{69} \pi (1 + \alpha)^2}. \tag{F.33.5}$$

Then $|\Gamma| \leq \min\{N, R_\alpha\}$. There are $Y \subseteq X$ and a Freiman homomorphism

$$\Psi : C \longrightarrow H, \quad C = \text{Bohr}(\Gamma, \zeta_\alpha),$$

such that

$$|Y| \geq \frac{\alpha p(\alpha)}{20} |X|, \quad \Psi(0) = 0, \tag{F.33.6}$$

and, whenever $y, z \in Y$ and $r(y) - r(z) \in C$,

$$\phi(y) - \phi(z) = \Psi(r(y) - r(z)). \quad (\text{F.33.7})$$

In fact Ψ can be taken to preserve every order-seven relation in C , with the same parameters.

The set Y can depend on ϕ , while Γ depends only on the fibre-count function g . This distinction matters in the multilinear stages of the source argument. No assumption of prime order is needed in theorem F.33.1; primality enters only if one wants to write a local additive map in cyclic coordinates as multiplication by a scalar.

F.33.2 What is improved relative to the inspected sources

The repository snapshot [131] includes the original paper’s transcription, the formal development, a compiled collection of five earlier reports, and supporting files for later reports. We additionally retrieved and read the historical manuscripts of source 07 [133] and source 12 [134]. Their earlier ideas are important inputs to the present development.

Table F.33.1: Local asymptotic controls as $\alpha \downarrow 0$. The constants are explicit in the cited statements. The volume exponent is L in a lower bound $N \exp(-O(L))$.

Quantity	Corrected statement	Section 10	Prior Bohr report, source 12	Present theorem
Retained fraction	$\alpha^6/20000$		$\alpha^5/20000$	$\alpha p(\alpha)/20$
Frequency count	$O(\alpha^{-10})$		$O(\alpha^{-10})$	$O(\alpha^{-3})$
Inner radius	$\exp[-O(\alpha^{-10} \log(1/\alpha))]$		$\Omega(\alpha^{19})$	$\Omega(\alpha^5)$
Volume exponent	$\alpha^{-20} \log(1/\alpha)$		$\alpha^{-10} \log(1/\alpha)$	$\alpha^{-3} \log(1/\alpha)$
Input order-eight defect	2^{-43}		2^{-43}	2^{-21}

The α^5 retention was already identified by source 12 and was available internally in the earlier formal argument. We do not claim it as a new result. Likewise, polynomial Bohr radii, the full-fourth-moment smoothing estimate, and coupled-path rigidity were already developed in sources 12 and 07. The new work is the logarithmic component lemma, the energy-weighted boundary estimate, their connection to a directly normalized anchor law, and the complete resulting theorem. The historical printed and corrected statements have different constants and density qualifications; the comparison in table F.33.1 uses the corrected repository statement and the explicit sparse branch of source 12, not the unrepaired printed density claim.

[write] **6 October 2026; sources 28–33.** *Identifications within this report.* Source 28 numbers the earlier manuscripts as this report does: [132] is source 04 (the base), [133] is source 07 (its chapter is Section II.11) and [134] is source 12 (Section II.33); “the earlier collection” and “a compiled collection of five earlier reports” are this report’s first write (sources 01–05), which was `article.tex` at the pin `e3125e698`. The first column of the table is the corrected formal statement `theorem_10_13` (`Section10.lean`, line 473; companion `theorem_10_13_holds`, `Proofs10Main.lean`, line 1292, at HEAD `5d4718772`): threshold $\lambda = 2^{-37}\alpha^{11/2}$, input defect 2^{-43} , $\alpha \leq 1/6$, retained $\alpha^6|X|/20000$ and spectrum at most $2^{74}\alpha^{-10}$, as printed (checked at the write). The second column is source 12’s Corollary II.40.3; its α^5 retention was already noted inside the formal proof (note in Section F.17). At $\alpha = 1/6$ the present $\alpha p(\alpha)/20$ is $1.08 \cdot 10^{-4}$ against $1.07 \cdot 10^{-9}$ for $\alpha^6/20000$ (computed at the write).

F.33.3 A second result: the higher-norm extremizer changes

For real centered functions on odd-order groups, the earlier collection proves

$$\|f\|_{U^3}^8 \geq 2\|f\|_{U^2}^8,$$

with equality for one conjugate pair of Fourier frequencies [132]. It asks for the sharp comparison at higher orders. It is therefore natural to test the same proposed extremizer at U^4 . This test fails.

If a character χ has order greater than 24, put

$$f_b = \chi + \chi^{-1} + b(\chi^3 + \chi^{-3}), \quad b \in \mathbb{R}.$$

The exact calculation in the second part of the article gives

$$\left. \frac{d}{db} \frac{\|f_b\|_{U^4}^{16}}{\|f_b\|_{U^2}^{16}} \right|_{b=0} = 54 > 0.$$

Consequently a sufficiently small *negative* third harmonic improves on the cosine. We prove an analogous nonstationarity statement for every $d \geq 4$, with an explicit sufficient character-order condition. This rules out a natural extrapolation; it does not determine the true optimal constants.

F.33.4 Status and reading guide

The mathematical claims are supported by complete arguments below. Some finite polynomial coefficients additionally have short, exact, reproducible enumerations. The manuscript has received independent derivation and adversarial checking within this work, but has not been peer reviewed or verified in Lean. The package records the checks rather than treating computation as a substitute for proof.

There is no claim of priority over the full literature. Modern inverse theorems, such as the quasipolynomial work of Leng, Sah, and Sawhney [3], address a broader quantitative theory. Our explicit local model is intended for the particular information-preserving interface in this repository. Improving that interface does not by itself close the later Section 13, Section 16, and Section 18 obligations.

F.34 Source 29. Purpose, contributions, and their precise scope

Source 29, delivered Section 1.

Added in the sixth write, 6 October 2026. Delivered title page. “ProveIt research manuscript / Further Quantitative Refinements of Gowers’s Szemerédi Argument / Higher norm comparisons, joint quadratic rigidity, and restriction by affine rank / Prepared for Vladimir Reshetnikov / Mathematical development and verification with ChatGPT / 6 October 2026”; *the PDF author field reads “ChatGPT; prepared for Vladimir Reshetnikov”*. Archive `gowers_moments_stability_restriction.zip` (arrival 624d7a42a), manuscript `gowers_moments_stability_restriction.tex` (2,256 lines, 30 pages), pin e3125e698 (its `data/source_audit.json`, shipped as `data/29-moment-stability-source_audit.json`). Printed below: the abstract, the closing lines of its title page and Section 1; Sections 2–4 form a chapter of Part III (Section III.66), Sections 5–8 are in Part II with source 31 (Section II.104), Section 9 is in Part II (Section II.119), Sections 10–12 are in Part V. Its macro `\code` (verbatim file names) is renamed `\codetwenty-nine`; its ‘question’ environment (“Question”, on the shared counter) is set in this report’s ‘Question’ style.

Abstract of source 29 (as delivered). We prove three families of local refinements connected to Gowers’s proof of Szemerédi’s theorem and to the existing ProveIt research reports. First, a positive Fourier-moment argument gives $\|f\|_{U^d}^{2^d} \geq 2^{-m} \binom{2^d}{m} \|f\|_{U^2}^{2^d}$, where $m = 2^{d-2}$, for real functions whose Fourier transform vanishes on the characters of order at most two. For centered functions on odd-order groups this improves the previously used monotonicity bound for every $d \geq 4$; in dimension four the coefficient improves from 4 to $35/8$. We give the full torsion-dependent moment polynomial, moment-level equality and stability, and bounds on the unresolved sharp higher-norm constants.

Second, we prove an explicit joint stability theorem at the absolute endpoint of the derivative-spectrum inequality. If a nonnegative selector with row sums at most one captures $1 - \varepsilon$ of the maximum squared derivative correlation, with $0 \leq \varepsilon \leq 2^{-9}$, then the function and selector are simultaneously within 24ε , in normalized squared L^2 error and missing graph mass respectively, of a quadratic phase and its derivative graph. A mixed version recovers a common quadratic part and a relative character for two functions. The proof works on every finite abelian group, including even-order groups, and the linear order of the error is optimal.

Third, we retain every affine-rank stratum in a previously introduced random graph restriction. For a fixed finite field and a fixed translation-invariant linear system with s free coordinates, the size threshold becomes proportional to $(\alpha\eta)^{-1}$, with retention proportional to $\alpha^{s+1}\eta^s|B|^s$. Repeated entries and degenerate tuples are included exactly. A weaker coefficient removes the threshold altogether. Complete proofs, finite checks, source comparisons, and further research questions are supplied.

Unrefereed research manuscript. No new Lean formalization is claimed. No improvement of the global Szemerédi bound is claimed.

The aim is to strengthen interfaces inside an inverse argument while preserving the hypotheses under which they can be used. Three kinds of information are retained here: the positive Fourier coefficients of an autocorrelation, the orthogonal residual of a derivative spectrum, and the actual affine rank of a configuration. Discarding these quantities too early loses constants, stability rates, or size thresholds.

The original reference is Gowers’s 2001 paper [1]. Its Lemma 3.10 bounds cube counts from uniformity, Proposition 6.1 extracts additive structure from derivative correlations, and Lemma 9.3 improves the proportion of relations respected after restriction. Our results concern these local themes. The normalizations and domains below are stated independently, so they can be compared with the corrected repository interfaces without treating different hypotheses as interchangeable.

F.34.1 What is added to the inspected material

The repository was inspected at the commit

e3125e6982abe2d44bd7cb9369063d4e69d41a2d.

The existing combined report [135] proves the sharp real-centered U^3/U^2 gap and an exact full-graph endpoint, and asks about higher norm gaps and joint stability. The restriction manuscript numbered source 17 [136] already introduces random linear maps on augmented graph points and asks whether affine-rank stratification can improve its exceptional-tuple threshold. These are the immediate mathematical starting points.

[write] **6 October 2026; sources 28–33.** *Identifications within this report.* [135] is this report at its first write (sources 01–05): the sharp U^3/U^2 gap is source 04’s Theorem III.3.4, the full-graph endpoint

source 02’s (Section II.4), and the questions are source 03’s Research question V.14.1 with source 04’s Question 3 (item III.1) and source 02’s Subsection 11.2 (item II.2). [136] is source 17; its random linear maps on augmented graph points are in its Section 7 (Section II.63), and its “second further question” is its Question 2 (Section V.71). [137] are sources 08 and 09. Source 28, written from the same pin, proves overlapping bounds for c_d^{odd} ; neither cites the other.

Topic	Previously available interface	Result proved here
Higher norms	The sharp U^3 gap propagated by monotonicity	Central-binomial coefficients improving every dimension $d \geq 4$, plus a torsion profile
Joint endpoint	Exact quadratic rigidity and intermediate alignment estimates	Explicit linear error bounds for both functions and selectors, with a mixed version
Finite-field restriction	Generic tuples retain their rank; all nongeneric tuples are charged probability one	Exact survival for every tuple, a linear inverse-accuracy threshold, and a threshold-free corollary

Table F.34.1: The comparisons concern the inspected repository statements. The underlying Fourier, majority-repair, and random-map methods are classical or already present in that repository.

The central-binomial estimate is an application of the classical multinomial moment method; related Steinhaus expansions appear, for example, in [121, Section 8.3]. Qualitative stability of near-maximal Gowers norms is established by Eisner and Tao [7, Theorems 1.1 and 1.4]. They also explain that their methods yield effective polynomial rates. Our contribution in that direction is the explicit joint derivative-selector interface, including unbounded L^2 functions and the mixed setting. The homomorphism repair used in its proof belongs to the self-correction tradition of Blum, Luby, and Rubinfeld [49], with the test-error comparison also familiar from [122]; we prove the precise version needed. The derivative correlations are finite short-time Fourier transforms. The quadratic geometry of their exact support minimizers is established in Nicola’s work [123]; our statements quantify a particular near-endpoint problem with a prescribed selector.

These are proved local extensions and answers to specified questions in the repository. A targeted literature check is not an exhaustive priority determination. In particular, the paper does not announce a new qualitative inverse theorem, a sharp value of the higher norm constants, or a solution of the unrestricted weighted relative-energy stability problem. The general bound of Leng, Sah, and Sawhney [2],

$$r_k(N) \ll_k N \exp(-(\log \log N)^{c_k}), \quad k \geq 5,$$

is a literature benchmark for comparison; no substitution through an entire density-increment iteration is performed here.

F.34.2 Guide to the manuscript

Sections III.67–III.69 develop the higher-norm bounds and their precise sharpness limitations. Sections II.110–II.112 establish joint quadratic rigidity and its consequences. Section II.120 proves the restriction results, including the generic higher-error-rank improvement. The final sections explain integration, verification, and open research directions. Every theorem in the three main families is proved in this manuscript; the included code supplies finite checks of identities and examples rather than substitutes for those proofs.

F.35 Source 30. Target, contributions, and comparison standard

Source 30, delivered Section 1.

Added in the sixth write, 6 October 2026. Delivered title block. “Fixed-radius Bohr invariance in Gowers’s Szemerédi argument / Escape-tube bounds, larger Freiman-extension neighborhoods, and exponential-rank obstructions / Research report prepared for Vladimir Reshetnikov / Mathematical development and exact checks with ChatGPT / 6 October 2026”; *the PDF author field reads “Research report prepared with ChatGPT”*. Archive `Gowers_Fixed_Radius_Bohr_Invariance.zip` (arrival 12da35673), manuscript `article.tex` (1,278 lines, 19 pages), pin 8d65d3c9f (its `SOURCE_AUDIT.md`, shipped as `30-fixed-radius-bohr-SOURCE_AUDIT.md`). Printed below: the abstract, the integration target and status lines after it, and Section 1; Sections 2–9 are in Part II (Section II.121), Sections 10–12 and Appendix A in Part V. Its macros `\defect` (an operator, not this report’s ε) and `\pin` are renamed `\defectthirty` and `\pinthirty`; its ‘question’ environment (“Question”, on its own counter) is set in this report’s ‘Question’ style, and its Questions 1–12 keep their indices. It numbers equations consecutively through the manuscript; here they are numbered within sections.

Abstract of source 30 (as delivered). We give a direct fixed-radius replacement for the Bohr translation estimate in Lemma 10.10 of Gowers’s proof of Szemerédi’s theorem. For a rectangular Bohr set B defined by r characters, with widths at most $1/4$, a shift whose coordinates are at most $1/(2m)$ of the corresponding widths satisfies

$$2m(|B| - |B \cap (B + d)|) \leq |B(3\rho/2)| - |B| \leq (3^r - 1)|B|.$$

The proof packs two families of disjoint outward translates of the escape sets into a larger Bohr shell. It uses neither a regular-radius selection nor a lower bound for the density of B in the ambient group. The first inequality is exactly sharp already for intervals in cyclic groups.

As a consequence, the corrected radius $\rho^r/(2^{r+4}r)$ in Corollary 10.11 can, for $0 < \rho \leq 1/4$, be replaced by $\rho/[8(3^r - 1)]$ while preserving its overlap margins, or by $\rho/(6 \cdot 3^r - 4)$ for the extension conclusion alone. We give a higher-order extension theorem, an explicit density-sensitive version of the compatibility step underlying Lemma 10.12, and exact finite checks. An explicit Chinese-remainder construction shows that exponential dependence on the number of characters is unavoidable in a universal fixed-radius translation bound, even for cyclic groups of composite order. The upper and lower exponential bases do not match. These are local refinements with complete mathematical proofs; no new global bound for Szemerédi’s theorem, literature-priority claim, or proof-assistant verification is asserted.

Integration target. `Combinatorics/Ramsey/Research/GowersSzemeredi/`

Status. Unrefereed research manuscript. Exact checks supplement, and do not replace, the proofs. The existing repository is not modified by this package.

F.35.1 The local loss being removed

Let $G = \mathbb{Z}/N\mathbb{Z}$ and let K be a set of r frequencies. We use the angular normalization

$$B(K, \rho) = \{x \in G : \|kx/N\|_{\mathbb{T}} \leq \rho \text{ for every } k \in K\}, \quad \mathbb{T} = \mathbb{R}/\mathbb{Z}.$$

Gowers’s Lemma 10.10 gives the estimate

$$|B \cap (B + d)| \geq (1 - 2^{r+1}r\rho^{-r}\zeta)|B|, \quad d \in B(K, \zeta). \quad (\text{F.35.1})$$

Its proof first bounds the number of points near coordinate boundaries in terms of the ambient cardinality N , and then compares N with $|B|$. The latter comparison introduces a power ρ^{-r} . This

is the specific loss targeted here, not the much larger collection of losses elsewhere in the Szemerédi argument [1, Section 10].

The repository’s corrected Corollary 10.11 uses

$$\zeta_{\text{old}} = \frac{\rho^r}{2^{r+4r}}. \quad (\text{F.35.2})$$

It explicitly divides by r . The printed paper instead multiplies by r ; that expression does not give the overlap claimed in the proof. Our comparison is with the corrected repository expression, not with the misprint. The exact declarations inspected are `lemma_10_10` and `corollary_10_11` in `Section10.lean` at the pin recorded in Section V.118 [138].

F.35.2 What is proved here

Write $c_r = 3^r - 1$. The principal conclusions are as follows.

- (1) **An intrinsic escape bound.** If $B = B(\chi, \rho)$, $0 < \rho_i \leq 1/4$, and $\|\chi_i(d)\|_{\mathbb{T}} \leq \rho_i/(2m)$, then

$$2m \operatorname{def}_B(d) \leq |B(\rho + m\mathbf{v})| - |B| \leq |B(3\rho/2)| - |B|, \quad v_i = \|\chi_i(d)\|_{\mathbb{T}}.$$

Here $\operatorname{def}_B(d) = |B| - |B \cap (B + d)|$. This holds in every finite abelian group, for unequal widths, without changing B .

- (2) **A larger extension neighborhood.** At density $7/8$ inside B , the radius $\rho/[8c_r]$ preserves the original $7/8$, $5/8$, and $1/8$ overlap margins. The still larger radius $\rho/(6c_r + 2)$ suffices for the Freiman extension conclusion itself.
- (3) **A rank obstruction.** For $r = 2s + 1$, a universal estimate of $\operatorname{def}_B(d)/|B|$ by a coefficient depending only on r times the relative shift must have coefficient at least $3^s/2$ in the small-shift limit. The examples can be taken in cyclic groups, but their orders are composite.

The auxiliary extension and compatibility arguments retain their exact hole and error budgets. This yields a general order- k statement and, in the setting spelled out in Section II.133, the sufficient condition $\sqrt{\theta} < 5/14$ instead of the coarser $5/16$ obtained by measuring both exceptional direction sets against $|B|$.

F.35.3 Relation to regular and smoothed Bohr sets

Regular-radius arguments and smoothing are established tools, not new ideas of this manuscript. Green explains the radial boundary problem and Bourgain’s regular-radius remedy, and develops smoothed Bohr neighborhoods [124, Sections 3–4]. The distinction here is that the input set, its radius, the dense subset, and the partial map are all kept unchanged. Changing the radius after a partial map has been constructed may change its domain and its relative density. The present argument pays an exponential cost in the number of characters to avoid that operation.

In particular, our result is not a radial regularity theorem. A Bohr set can have many points exactly on its boundary, so even a tiny radial contraction can remove many points. A translation by an actually available small group element is a different operation. The proof below exploits that element’s orbit and the long interval of non-return after a boundary exit.

F.35.4 Scope and novelty discipline

All numbered assertions labeled theorem, lemma, proposition, or corollary below have proofs. Questions are stated separately. The new research focus is the fixed-radius escape inequality and its quantitative consequences; the elementary packing and Freiman telescoping arguments are presented as supporting tools, not as newly discovered principles. The literature check is targeted, not exhaustive. The explicit family in Section II.135 refutes a precisely formulated possible polynomial-rank improvement; we do not identify that possibility as a previously published open conjecture.

The work does not change the formal status of any repository result. It does not establish a new bound for the largest progression-free subset of $\{1, \dots, N\}$, and does not assert an improved full Theorem 10.13 merely from substitution into one local step.

[write] **6 October 2026; sources 28–33.** *Credit within this report.* Source 30 does not cite source 12 (its manuscript was in the repository’s history and its code in the tree at the pin), and two of its supporting steps are source 12’s: the tree budget (Lemma II.131.2) is source 12’s Lemma II.34.3, and the compatibility inequality of Theorem II.133.1 is that of source 12’s Lemma II.39.1. Its extension budget is weaker than source 12’s sharp Theorem II.35.2 (notes there). What is new here is the fixed-radius escape inequality (Theorems II.129.3, II.130.2) with its sharp constituents, the radius comparison with the corrected Corollary 10.11 at a fixed radius, the 5/14 count against the formal 5/16, and the exponential-rank obstruction in cyclic groups (Section II.135). Source 12’s Corollary II.37.3 has the much larger radius $\rho/(8d)$, but at a radius selected in $[\rho/2, \rho]$; the obstruction here explains why a fixed-radius statement must pay exponentially in the rank.

F.36 Source 31. The problem, the contributions, and their scope

Source 31, delivered Section 1.

Added in the sixth write, 6 October 2026. Delivered title block. “Weighted Derivative Spectra and Sharp Local Rigidity / Explicit Freiman restriction, an optimal near-extremal U^3 profile, and the limits of stability for correlation weights / Prepared for Vladimir Reshetnikov / Mathematical development and verification with ChatGPT / 6 October 2026”; *the PDF author field reads “Research developed with ChatGPT for Vladimir Reshetnikov”*. Archive `gowers_weighted_spectrum_rigidity.zip` (arrival 12da35673), manuscript `gowers_weighted_spectrum_rigidity.tex` (1,745 lines, 25 pages), pin `e3125e698` (its `SOURCE_LEDGER.json`, shipped as `data/31-weighted-spectra-SOURCE_LEDGER.json`). Printed below: the abstract and Section 1; Sections 2–5 and 9 are in Part II (Section II.104), Sections 6–8 and Appendix A in Part II (Section II.114), Sections 10–11 and Appendix B in Part V. Its macros `\Energy` (a sans-serif E) and `\ip` are renamed `\Energythirtyone` and `\ipthirtyone`; its ‘question’ environment (“Research question”, on the shared counter) is set in this report’s ‘Research question’ style.

Abstract of source 31 (as delivered). We develop three connected refinements at the Fourier-to-structure interface of Gowers’ proof of Szemerédi’s theorem. The first is an exact local extremal result. For every finite abelian group and every nonzero complex function f , put $\kappa = \|f\|_{U^3}^8 / \|f\|_2^8$. If $\delta = 1 - \kappa < 1/120$, then the best squared correlation with a normalized quadratic phase supported on a subgroup coset is at least

$$1 - \mathcal{R}(\delta), \quad \mathcal{R}(\delta) = \frac{1 - \sqrt{(1 + \sqrt{1 - 4\delta})/2}}{2}.$$

This function is optimal uniformly over finite abelian groups, with equality already on the two-element group. A spectral subgroup argument gives an explicit preliminary bound independent of the group exponent; a quartic energy calculation then gives the exact profile. We obtain

a sharper full-graph endpoint: a derivative graph of score at least $1 - \varepsilon$ admits a symmetric homomorphism with at most ε errors, and quadratic-phase infidelity at most $(1 - \sqrt{1 - 2\varepsilon})/2$, for $\varepsilon \leq 1/400$.

Second, we carry prescribed nonnegative correlation weights through additive-energy extraction and an exact Freiman restriction. A weighted version of the prime-field restriction has retention $1/(kC^{2k+1})$, compared with $1/(8kC^{4k})$ in Gowers’ Lemma 7.5. The underlying smaller sumset argument is already present in Green–Tao; we give its finite weighted form and an explicit composition retaining the original derivative correlation. An order-eight specialization improves the historical Corollary 7.6 parameters to $2^{-363}c^{171}$ by an elementary route, with a separate modern comparison using Reiher–Schoen.

Third, we prove a defect inequality that controls Fourier uniformity on the extremal subgroup and the squared masses of other cosets. An exact total-variation identity gives linear transfer to a nearby input. Explicit examples show why near equality cannot control the weights in normalized L^1 without additional complexity assumptions. All principal refinements have proofs. The article distinguishes these results from classical inverse theorems, gives reproducible finite checks, and proposes further research. It does not claim a new global Szemerédi bound or Lean certification.

Gowers’ argument turns information about Fourier coefficients of derivatives into additive structure of the frequency labels. Several losses can occur between these two descriptions: selecting one frequency from a large fiber, replacing a weight by the size of its support, extracting a small-doubling subset, and restricting further until every relevant additive relation is respected. Near an extremal configuration, a different question becomes available: can the entire function and its selected frequencies be recovered with an explicit error bound?

This article treats both regimes. For arbitrary positive correlation it gives weighted extraction with explicit, though deliberately conservative, polynomial losses. Near the maximum it identifies an exact optimal stability profile. The same derivative spectrum connects the two arguments.

F.36.1 The repository checkpoint

The source checkpoint is the following revision of `VladimirReshetnikov/ProveIt`:

e3125e6982abe2d44bd7cb9369063d4e69d41a2d.

We inspected the Ramsey transcription of Gowers’ paper, the consolidated *Local Quantitative Refinements* report, the relevant Lean declarations, and the three newly arrived reports on characteristic five, degree stratification, and sharp thresholds [139, 140, 141]. The present subject is distinct from their centered cube-count expansions, polynomial-coloring thresholds, and arrangement-degeneracy estimates.

Two questions in the consolidated report are especially relevant. Its Theme II asks for a structural extraction that retains specified derivative correlation, and for a description of near equality in the coupled function–weight energy inequality. The weighted restriction below answers a prime-cyclic bounded-fiber version of the first question. The sharp phase theorem and the weight counterexamples provide a positive theorem and a necessary limitation for the second. These are precise local answers; compatibility across the later multilinear slices remains a separate issue.

[write] 6 October 2026; sources 28–33. *Identifications within this report.* The “consolidated report” [140] is this report’s first write (sources 01–05); its weighted derivative energy and full-graph rigidity are source 02’s (Sections II.2–II.4), its elementary BSG source 03’s (Section II.5), and the two Theme II questions

are items II.1 and II.2. [141] are sources 21–23. Not seen at the pin: source 10’s Sections II.20–II.21 and source 11’s chapter (Section II.22), whose texts were not in this report’s article at the pin, bear on Sections II.108–II.117 (notes there); source 29, written from the same pin, treats the endpoint with weaker constants.

F.36.2 A map of the results

Result	Conclusion	Range
Theorem II.107.4	Exact optimal near-extremal infidelity $\mathcal{R}(1 - \kappa)$ to a quadratic coset phase	Every finite abelian group; $1 - \kappa < 1/120$
Theorem II.107.3	Quartic upper envelope, a quantitative deficit, and a stricter odd-order envelope	Fidelity $1 - t$, $t \leq 1/6$
Theorem II.108.1	Sharp separate phase and graph error bounds, with a joint error inequality	Full or nearly full graph; $\varepsilon \leq 1/400$
Theorem II.116.1	Weight retention $1/(kC^{2k+1})$ under small graph difference set	Prime target field; $k \geq 2$
Theorem II.116.2	Prescribed weight retained on an exact order- k Freiman graph	Bounded fibers, positive weighted energy
Theorem II.117.1	Original squared derivative correlation survives the complete restriction	Prime-cyclic domain; explicit spectrum profile
Theorem II.109.1	Defect controls U^2 oscillation and coset-mass leakage; L^1 under complexity bounds	Any extremal subgroup
Appendix II.118	Weight mass at least $(1 - \epsilon)\sqrt{c}M$ with polynomial doubling	Consequence of Reiher–Schoen

F.36.3 Relation to established results

The qualitative critical-norm statement is established mathematics. Eisner–Tao [7, Theorem 1.4] classify extremizers and near-extremizers of the critical Gowers norm inequality on locally compact abelian groups. At order three the critical Lebesgue exponent is two, and the extremizers are precisely normalized quadratic phases on cosets. We reprove the finite case needed here and determine an explicit sharp local modulus.

The phase-space language also overlaps directly with stabilizer-state testing. Gross–Nezami–Walter [125, Sections 2–3] use Weyl characteristic functions, isotropic subspaces, and large-coefficient arguments. Our finite-group argument first forces a subgroup by a small-difference-set lemma and then proves isotropy by averaging operators. This avoids a lower bound for the separation between distinct roots of unity. The resulting constants do not depend on the exponent of the group. The exact local quartic envelope is proved below. A bounded literature search did not establish an earlier occurrence of this exact profile; independent bibliographic priority is not certified.

Weighted Balog–Szemerédi–Gowers principles are also classical; see [126, 127]. The smaller vertical-defect sumset used in Section II.116 already appears in Green–Tao [4, proof of Lemma 6.2]. We credit that mechanism explicitly. The contributions at this interface are the weighted finite

statement, the parameter ledger, and the preservation of a specified correlation through the entire local extraction. The modern set theorem of Reiher–Schoen [12] is stronger than the elementary extraction used for our self-contained numerical comparison; both routes are recorded.

Meaning of sharpness. The modulus in Theorem II.107.4 is optimal over the class of all finite abelian groups. The two bounds in Theorem II.108.1 are separately optimal. We do not claim that its joint error inequality, the cutoff $1/120$, the odd-order quartic envelope, or the extraction exponents are optimal. A result proved here is not thereby asserted to be a new global inverse theorem or a solution of an unrestricted open problem.

F.37 Source 32. The theorem and its scope

Source 32, delivered Section 1.

Added in the sixth write, 6 October 2026. Delivered title block. “Threshold free quadratic density increments on cyclic groups / Report276 / Research report prepared for Vladimir / 6 October 2026”; the PDF author field reads “Research report prepared for Vladimir”. Archive `Report276_Threshold_Free_Quadratic_Density_Increments.zip` (arrival `6c2e2172b`), manuscript `article.tex` (533 lines, 12 pages), pin `c94998a5d` (its `SOURCES.md`, shipped as `32-quadratic-increments-SOURCES.md`). Printed below: the abstract and Section 1; Sections 2–5 continue source 26’s chapter of Part I (Section I.40), Sections 6–7 are in Part V. It numbers equations consecutively through the manuscript; here they are numbered within sections.

Abstract of source 32 (as delivered). Let $A \subseteq \mathbb{Z}/N\mathbb{Z}$ have density δ . Suppose that M disjoint proper modular arithmetic progressions cover A , their lengths are pairwise within a factor of two, and degree at most two polynomial phases have summed balanced correlation at least αN , where $\alpha > 0$. We prove, for every cyclic modulus and with no additional scale hypothesis, that a proper progression Q satisfies

$$|Q| \geq \max \left\{ 1, \frac{(N/M)^{1/39}}{17} \right\}, \quad \frac{|A \cap Q|}{|Q|} \geq \delta + \frac{\alpha}{3}.$$

The proof includes an explicit quadratic recurrence estimate, a lossless phase-flat partition, and an all-moduli second-moment argument that handles the small-correlation regime. An explicit singleton comparison shows that the length guarantee also dominates $(N/M)^{1/64}/8$, hence the whole degree-two specialization of a pinned repository Corollary 5.8. All constants and rounding choices are included. The recurrence and localization mechanisms are classical; the report makes no literature-priority, formal-verification, or global Ramsey or Szemerédi bound claim. Its executable companion checks bounded exact identities and regressions, rather than certifying the real-parameter theorem numerically.

Write $G = \mathbb{Z}/N\mathbb{Z}$, where $N \geq 2$ is an integer. For $t \in \mathbb{R}$ let

$$e(t) = \exp(2\pi it), \quad \|t\| = \min_{m \in \mathbb{Z}} |t - m|.$$

For $y \in G$, define $e_N(y) = e(\tilde{y}/N)$ using any integer representative $\tilde{y}$. A *proper modular arithmetic progression* is a set

$$P = \{u + jd : 0 \leq j < n\} \subseteq G$$

whose displayed parameterization is injective; its length is n . A nonzero step need not be invertible. Singletons are proper, and progressions are allowed to wrap in their standard integer representatives.

Let $A \subseteq G$, $\delta = |A|/N$, and $f = \mathbf{1}_A - \delta$. The phases in this report are ordinary coefficient polynomials over G :

$$\phi_i(x) = a_i x^2 + b_i x + c_i, \quad a_i, b_i, c_i \in G.$$

Thus degree at most two includes affine and constant phases. No broader definition by vanishing third differences is assumed over composite moduli.

Theorem F.37.1 (Quadratic local cover increment). *Let $M \geq 1$ and let $P_1, \dots, P_M$ be nonempty, pairwise disjoint, proper modular progressions. Suppose*

$$A \subseteq U := \bigcup_{i=1}^M P_i, \quad |P_i| \leq 2|P_j| \quad (1 \leq i, j \leq M). \quad (\text{F.37.1})$$

Let ϕ_i have the form above, and suppose $\alpha > 0$ and

$$\sum_{i=1}^M \left| \sum_{x \in P_i} f(x) e_N(-\phi_i(x)) \right| \geq \alpha N. \quad (\text{F.37.2})$$

There is a nonempty proper modular progression Q such that

$$|Q| \geq \max \left\{ 1, \frac{(N/M)^{1/39}}{17} \right\}, \quad |A \cap Q| \geq \left(\delta + \frac{\alpha}{3} \right) |Q|. \quad (\text{F.37.3})$$

There is no lower bound on α or N/M , and no primality or least-prime-factor condition on N . The progression Q need not be contained in U .

[write] 6 October 2026; sources 28–33. *Against the open formal statement* (checked at the write at HEAD 5d4718772). `corollary_5_8` (Section05.lean, line 242) quantifies over every degree $k \geq 1$; with $k = 2$, `PolynomialOn 2` (Definitions.lean, line 222) is the coefficient-polynomial convention here and `polynomialPartitionConstant 2 = (2!)^2 2^9 = 2048` (Section05.lean, line 45), so the degree-two instance asks for gain $\alpha/16$ at length $(N/M)^{1/2048}/8$. By (V.125.1) this theorem implies that instance at every scale (Section V.125); with source 26’s Corollary I.41.3 for $k = 1$, the instances $k \leq 2$ have written proofs. The ledger is unchanged, because the statement quantifies over every degree; a formalization candidate (the natural target is a separate theorem for $k \leq 2$), not formalized. The bound exceeds one only when $N/M > 17^{39} \approx 9.7 \cdot 10^{47}$; below that the conclusion is a singleton, as the source says. It does not give the $k = 1$ instance, whose length $(N/M)^{1/16}/8$ grows faster.

[write] 7 October 2026; sources 40–46. *Every degree* (source 41, Section I.62). Source 41’s Theorem F.46.1 gives gain $\alpha/3$ at length $\max\{1, (N/M)^{1/D_k}/8\}$ in every degree k ; at $k = 2$ it keeps this theorem as its Corollary I.68.5, which is sharper there ($1/39$ against $1/513$). With sources 26 and 32, all instances of the formal `corollary_5_8` have written proofs; the formal statement is now exact as well, by the repository’s own proof (`corollary_5_8_holds`; note in Section F.46).

The proof uses two regimes. A second moment supplies the requested increment when α is small. Otherwise the correlation forces the cover to have enough mass that its shortest parent is long enough for an explicit quadratic phase partition. This avoids imposing a phase-partition threshold as an additional theorem hypothesis.

Section I.45 proves the uniform recurrence bound $p \leq 1024R^5$. Section I.46 turns it into a partition whenever the parent length is at least $2^{116}L^{38}$. Section I.47 proves the second-moment alternative over every cyclic group. Section I.48 chooses the target length and assembles the conclusion. The exact source comparison in Section V.125 is about a written degree-two implication; no repository file has been changed or Lean proof supplied.

F.38 Source 33. The problem and the contribution

Source 33, delivered Section 1.

Added in the sixth write, 6 October 2026. Delivered title block. “Sharp High-Agreement Stability for Gowers’s Respected Arrangements / Full-product decoding, optimal error terms, and erasures / Research report prepared for Vladimir Reshetnikov / Mathematical development and exact computational checks with ChatGPT / 6 October 2026”; *the PDF author field reads “Research report prepared with ChatGPT for Vladimir Reshetnikov”*. Archive `Gowers_Respected_Arrangement_Stability.zip` (arrival 17dd7886c), manuscript `article.tex` (1,623 lines, 24 pages), pin 9f83dbe0f (its `SOURCE_AUDIT.md`, shipped as `33-arrangement-stability-SOURCE_AUDIT.md`). Printed below: the abstract, the status statement after it and Section 1; Sections 2–8 are in Part II (Section II.136), Sections 9–10 and Appendix A in Part V. Its macros `\V` (the model class $\mathcal{V}$) and `\pin` are renamed `\Vthirtythree` and `\pinthirtythree`; its ‘question’ environment (“Research question”, on its own counter) is set in this report’s ‘Research question’ style, and its Research questions 1–12 keep their indices.

Abstract of source 33 (as delivered). Gowers’s proof of Szemerédi’s theorem uses arrangements of vertical edges whose horizontal coordinates obey one balanced additive relation. We study the high-agreement inverse problem for these same arrangements on a full product of finite abelian groups. The exact models are not just biaffine maps: they are $\phi(x, y) = a(x) + B(x, y) + \ell(y)$, with an arbitrary row function a , a biadditive map B , and a homomorphism ℓ .

For a test involving an even number $s \geq 4$ of vertical differences, write ε_s for its failure probability and δ for distance to this model class. We give a self-contained, dimension-free decoder and prove that, when $\varepsilon_s \leq 1/(240(s-1))$,

$$\delta \leq \frac{1 - \sqrt{1 - 8(s-1)\varepsilon_s/s}}{4(s-1)} = \frac{\varepsilon_s}{s} + \frac{2(s-1)}{s^2}\varepsilon_s^2 + O_s(\varepsilon_s^3).$$

Both displayed asymptotic coefficients are optimal uniformly over finite abelian groups. The leading coefficient remains optimal even for $G = H = K = \mathbb{F}_p$, with p an odd prime tending to infinity. For $s = 4$ we prove the stronger local inequality $\varepsilon_4 \geq 4\delta - 24\delta^2 + 32\delta^3$ for $\delta \leq 1/4$; an explicit family attains this entire cubic. For Gowers’s eight-arrangements, $s = 16$, the repair bound starts with $\varepsilon_{16}/16 + 15\varepsilon_{16}^2/128$. A further moment bound sharpens the high-density case of Gowers’s Lemma 12.3: θN^8 respected pairs imply at least $\max\{\theta^7, 4\theta - 3\}N^{32}$ respected eight-arrangements. We also retain a row-collision profile, treat erasures, give exact cyclic counting and decoding algorithms, and report reproducible integer-arithmetic checks. These are high-agreement full-domain results, not an improvement to the general sparse restriction step or to the global bounds for Szemerédi’s theorem.

Status. Self-contained research manuscript with proofs and finite computational checks. Unreferenced; not Lean-verified. The mathematical claims below are the claims proved in this manuscript, not a certification of literature priority. No human authorship or independent referee review is implied.

F.38.1 From many respected arrangements to a global model

In Section 12 of [1], Gowers first obtains many respected arrangements and then restricts the domain so that almost all arrangements are respected. The subsequent structural argument finds bilinear pieces. The present paper isolates a different endpoint question: when the domain is already a full product and the respected fraction is close to one, how many function values must be changed to obtain an exact model?

The distinction between a full and a sparse domain is essential. A full product supplies independent vertical base points, uniform marginals, and an exact derivative cocycle on every pair of increments.

On a sparse set these features need not survive conditioning. Our argument exploits these full-domain properties rather than silently assuming them in the sparse case.

Throughout, G, H, K are finite abelian groups, written additively. All probabilities are uniform on the finite sets specified, and all auxiliary random choices are independent unless a relation is explicitly imposed. We allow trivial groups. Let $s = 2m \geq 4$. A random test chooses $h \in H$, $x_1, \dots, x_s \in G$ uniformly subject to

$$x_1 + \dots + x_m = x_{m+1} + \dots + x_s, \quad (\text{F.38.1})$$

and independently chooses $y_1, \dots, y_s \in H$. It accepts $\phi : G \times H \rightarrow K$ when the corresponding vertical differences satisfy the same relation:

$$\sum_{i=1}^m D_h \phi(x_i, y_i) = \sum_{i=m+1}^s D_h \phi(x_i, y_i), \quad D_h \phi(x, y) = \phi(x, y+h) - \phi(x, y). \quad (\text{F.38.2})$$

This is Gowers's m -arrangement, with $s = 2m$ edge occurrences and $2s = 4m$ point occurrences. We count repetitions and the increment $h = 0$. Let $\varepsilon_s(\phi)$ be its failure probability.

Definition F.38.1 (The exact vertical model class). Let $\mathcal{V}(G, H; K)$ consist of the functions

$$\psi(x, y) = a(x) + B(x, y) + \ell(y), \quad (\text{F.38.3})$$

where $a : G \rightarrow K$ is arbitrary, $B : G \times H \rightarrow K$ is additive in each argument, and $\ell : H \rightarrow K$ is a homomorphism. Put

$$\delta(\phi) = \min_{\psi \in \mathcal{V}(G, H; K)} \mathbb{P}_{x, y}[\phi(x, y) \neq \psi(x, y)].$$

The class is a subgroup of the group of all functions. Its arbitrary $a(x)$ is necessary: vertical differences cannot see an added function of x . In particular, no test of the form (F.38.2) alone can force a to be affine. Every model passes every test, since $D_h \psi(x, y) = B(x, h) + \ell(h)$.

F.38.2 Main theorems

Theorem F.38.2 (Dimension-free high-agreement repair). *Let $s \geq 4$ be even, and write $\varepsilon = \varepsilon_s(\phi)$. If*

$$0 \leq \varepsilon \leq \frac{1}{240(s-1)}, \quad (\text{F.38.4})$$

then some $\psi \in \mathcal{V}(G, H; K)$ differs from ϕ on at most $R_s(\varepsilon)|G||H|$ inputs, where

$$R_s(u) = \frac{1 - \sqrt{1 - \frac{8(s-1)}{s}u}}{4(s-1)}. \quad (\text{F.38.5})$$

In particular,

$$\delta(\phi) \leq \frac{\varepsilon}{s} + \frac{2(s-1)}{s^2} \varepsilon^2 + O_s(\varepsilon^3). \quad (\text{F.38.6})$$

The leading coefficient $1/s$ is optimal even on common odd prime cyclic groups. With that leading coefficient fixed, the quadratic coefficient $2(s-1)/s^2$ cannot be decreased uniformly over all finite abelian groups.

[write] **6 October 2026; sources 28–33.** *In this report.* The first full-product repair theorem for Gowers’s arrangements, with sharp first and second coefficients; it advances item II.24 (source 11’s Question 7, source 14’s Questions 5–6, source 15’s Questions 2–3) on full or near-full products; the sparse case is open (item II.25). At $s = 16$: $1/16$ and $15/128$; the common-prime bracket of Research question V.127.1 is $23/512 \leq c_{16}^{\text{prime}} \leq 15/128$ (checked at the write).

[write] **7 October 2026; sources 60–69.** *Every order and every coefficient* (source 60, Section II.331): $\varepsilon_s \geq B_s(\delta)$ for $\delta \leq 1/4$ and every even $s \geq 4$, with all interior equality cases, and every Taylor coefficient of the repair bound optimal in succession (its Theorems II.334.1, II.334.3 and Corollary II.335.3).

The repair statement is proved in Section II.140; the two optimality statements are proved in Section II.141. The numerical threshold in (F.38.4) is a convenient explicit one and is not claimed optimal.

Theorem F.38.3 (A sharp cubic and a higher-order local bound). *If $\delta = \delta(\phi) \leq 1/4$, then for every even $s \geq 4$,*

$$\varepsilon_s(\phi) \geq F_s(\delta), \quad F_s(v) = sv(1-2v)((1-2v)^{s-2} - (2v)^{s-2}). \quad (\text{F.38.7})$$

In particular,

$$\varepsilon_4(\phi) \geq 4\delta - 24\delta^2 + 32\delta^3. \quad (\text{F.38.8})$$

For each odd $n \geq 3$, equality in (F.38.8) is attained with $G = \mathbb{Z}_n$, $H = K = \mathbb{Z}_2$, and $\delta = 1/(2n)$. For $s > 4$ we do not assert that the whole function F_s is optimal.

Corollary F.38.4 (Gowers’s eight-arrangements). *Suppose a function $\phi : G \times H \rightarrow K$ respects at least a fraction $1 - \varepsilon$ of its eight-arrangements and $\varepsilon \leq 1/3600$. Then*

$$\delta(\phi) \leq \frac{1 - \sqrt{1 - \frac{15}{2}\varepsilon}}{60} = \frac{\varepsilon}{16} + \frac{15}{128}\varepsilon^2 + O(\varepsilon^3). \quad (\text{F.38.9})$$

For $G = H = K = \mathbb{F}_p$, the approximating function has the form $a(x) + (bx + c)y$, where $b, c \in \mathbb{F}_p$ and a is arbitrary.

Proof. Take $s = 16$ in Theorem F.38.2. A biadditive map $\mathbb{F}_p \times \mathbb{F}_p \rightarrow \mathbb{F}_p$ is determined by its value at $(1, 1)$, and a homomorphism $\mathbb{F}_p \rightarrow \mathbb{F}_p$ by its value at 1. $\square$

F.38.3 Position relative to the source and existing refinements

The primary source is Gowers’s 2001 paper [1]; the repository’s corrected transcription [143] was used to check the arrangement definition and the nearby lemma interfaces. Table F.38.1 records the logical relationship rather than claiming that a theorem with stronger hypotheses replaces its sparse counterpart.

The existing local-refinements report in the repository treats, among other topics, density extraction, energy inequalities, moment amplification, transport rigidity and cubical expansions [143]. The elementary majority correction used here belongs to the self-testing tradition of Blum–Luby–Rubinfeld [49]; it is not presented as a new general principle. Our specific target is a row-gauge arrangement code and its quantitative defect-to-distance relation. The proof includes all required self-correction estimates, so no theorem from an unrefereed repository report is an unproved input.

[write] **6 October 2026; sources 28–33.** *Identifications within this report.* At this source’s pin the “existing local-refinements report” [143] was its second write (sources 01–10). The arrangement material

Source location	Role in the original argument	Present contribution
Definition after Lemma 12.4	Balanced relation among vertical differences	Same test on arbitrary finite abelian products; exact normalization.
Lemmas 12.2–12.4	Lower bounds for numbers of respected configurations	Near-one repair, and Lemma 12.3 supplemented by the bound $\max\{\theta^7, 4\theta - 3\}N^{32}$; the exponent seven is not lowered uniformly.
Lemmas 12.5–12.6	Selection of a subset with few violated arrangements	A repair theorem after full-product or near-full-product coverage is available.
Section 13	Bilinear structure on pieces of a sparse domain	Global biadditive structure modulo an unavoidable row function, under full-domain high agreement.

Table F.38.1: Source crosswalk. Here $s = 16$ means Gowers’s eight-arrangement.

it does not cite: source 14’s equality theory for arrangement moments (Section II.44) and source 15’s collision-sensitive amplification (Section II.45; its Lemma II.47.2 is the moment interpolation behind Gowers’s θ^7). The high-agreement repair asked for by source 11’s Question 7 and sources 14 and 15 (label repair near a product) is the subject here (item II.24).

For broader context, Gowers–Milićević study Freiman multi-homomorphisms on dense subsets of products of finite vector spaces [129]. Their sparse-domain structure problem is substantially more general in domain geometry and has different conclusions. Directional uniformity norms also provide a related setting [130]. Neither comparison constitutes a priority claim for the constants here.

There is no deduction here of a new upper bound for progression-free sets. For example, the global estimates of Leng–Sah–Sawhney [2] require machinery and parameter propagation not supplied by a full-domain high-agreement lemma alone. The proposed sparse extension is left explicitly as a research question.

F.39 Source 34. Scope, principal results, and the audited baseline

Source 34, delivered Section 1.

Added in the seventh write, 6 October 2026. Delivered title page. “Combinatorics and additive structure / Prime Torsion and Norm Gaps in Gowers Cube Counts / Minimal Boolean supports, the cost of positivity, and sharper uniformity inequalities / Research report prepared for Vladimir Reshetnikov / Mathematical development and exact computation with ChatGPT / 6 October 2026”; *the PDF author field reads “ChatGPT; research report prepared for Vladimir Reshetnikov”*. Archive `Gowers_Cube_Torsion_and_Norm_Gaps.zip` (arrival 70c6a1cf1), manuscript `gowers_cube_structure.tex` (2,805 lines, 37 pages), pin 12da35673 (its `source_ledger.json`, shipped as `data/34-torsion-gaps-source_ledger.json`). Printed below: the abstract of its title page (it has no abstract environment) and Section 1; Sections 2–9 form a chapter of Part III (Section III.73), Section 10 continues the chapter on the odd-order constants in Part III (Section III.71), Appendix A is in Part II (Section II.162), Sections 11–12 are in Part V. Its macro `\file` (small type) is renamed `\filethirtyfour`; it already uses a ‘problem’ environment printed “Research question” on the shared counter, as this report does.

Abstract of source 34 (as delivered). We develop local refinements of the cube-counting and rigidity steps in Gowers’s proof of Szemerédi’s theorem. Let $h(p)$ be the least support of an affinely independent Boolean configuration admitting a full-support characteristic- p relation. We prove $h(p) = D_{\text{div}}(p) + 1$, where $D_{\text{div}}(p)$ is the least order of a nonsingular binary matrix whose

determinant is divisible by p . This identifies the first dimension permitting any p -primary Smith torsion, and gives $h(7) = 7$ and $h(11) = h(13) = h(17) = 8$.

Retaining integral equivalence in Shitov’s composition, we prove that every positive integer up to $\exp(\frac{1}{2}N \log N - N \log \log N - 4N)$ is the sole nonunit Smith factor of a binary matrix of order N , eventually. The published Alon–Vu construction is the only deep input. Consequently, even exact cyclic quotients of order m have optimal support $(2 + o(1)) \log m / \log \log m$. We also characterize universal nonnegativity of cyclic Fourier moments: the relation coefficients must occur in opposite pairs. The corresponding optimal support is $(4 + o(1)) \log m / \log \log m$, so positivity doubles the leading cost.

An exact fourth dual-function calculation, combined with positive resonance terms, gives $\|f\|_{U^2}^4 \leq (2/3)\|f\|_{U^4}^4$ for every real centered function on every finite odd-order abelian group. A two-pair example on $\mathbb{Z}/29\mathbb{Z}$ gives ratio $0.558729\dots$, exceeding every single-pair ratio. This improves the leading fourth-dimensional cube coefficient to $200/3$. An appendix proves the sharp nonuniform directional rigidity threshold $[2 + (2k - 2)\kappa]^{-1}$. Full proofs, exact finite certificates, source attribution, and a research and formalization program are provided.

Manuscript status. Complete mathematical arguments and exact finite certificates are supplied, with internal proof review. External peer review and Lean certification remain separate steps. Statements of improvement refer to the sources identified below.

Gowers’s Lemmas 3.8–3.10 supply a multilinear Cauchy–Schwarz inequality, the triangle inequality for the uniformity norm, and an upper bound on cube counts [1]. These are natural places to ask how centering, Fourier symmetry, and the arithmetic of the ambient group improve a general-purpose norm estimate. The ProverIt research collection already develops fourth- and fifth-order cube expansions and explicitly asks for the first appearance of prime torsion and the sharp higher-dimensional norm gaps [155]. We concentrate on those questions.

The repository was inspected at commit

12da356730b12cb73c276a258aaadf3223e3662b.

The merged research article contains sources 01–10. Later numbered ancillary files and incoming archives are distinct parts of the audited material; they are not silently counted as chapters of that merged article. In particular, the incoming report *Further Quantitative Refinements of Gowers’s Szemerédi Argument* already proves the stronger baseline $c_4^{\text{odd}} \leq (8/35)^{1/4}$ [156]. The comparison in this report uses that bound, as well as the merged article’s earlier $2^{-1/2}$ bound.

F.39.1 What is established here

The structural parameter $h(p)$ is defined precisely in Section III.74; $h_{\text{cyc}}(m)$ imposes the stronger condition of an exact cyclic quotient. Our principal conclusions are:

1. **An exact extremal reduction.** For every prime p ,

$$h(p) = D_{\text{div}}(p) + 1.$$

The same value is the minimum number of Boolean vertices whose row quotient has any p -torsion. The minimum ambient cube dimension is $D_{\text{div}}(p)$. The result resolves the apparent extra full-support obstruction at the minimum dimension, while preserving that obstruction for a general nonminimal matrix.

2. **Optimal asymptotic support.** As p tends to infinity through primes and m through integers,

$$h(p) \sim h_{\text{cyc}}(p) \sim \frac{2 \log p}{\log \log p}, \quad h_{\text{cyc}}(m) \sim \frac{2 \log m}{\log \log m}.$$

This closes the leading-order gap between the repository's Hadamard lower bound and its linear upper bound $h(p) \leq p + 2$.

3. **An improved consecutive range with prescribed Smith form.** For every sufficiently large N , each integer

$$1 \leq m \leq \left\lfloor \exp \left(\frac{N}{2} \log N - N \log \log N - 4N \right) \right\rfloor$$

is realized by a binary matrix with Smith form $\text{diag}(1, \dots, 1, m)$. The construction follows Shitov's architecture [147]; the smaller secondary loss and the cyclic Smith-form conclusion are proved with all rounding and integral-equivalence details retained. A further parameter optimization allows any linear-loss constant greater than $2 \log 2 + \frac{1}{2} \log(3/(2 \log 2)) + \frac{1}{2} = 2.272283 \dots$ in place of 4.

4. **Exact small cases and an elementary detector.** We prove

$$h(7) = h_{\text{cyc}}(7) = 7, \quad h(p) = h_{\text{cyc}}(p) = 8 \quad (p = 11, 13, 17).$$

For every $m \geq 2$, a separate elementary construction gives a nonnegative Fourier moment on at most $2 \lceil \log_2 m \rceil + 2$ Boolean vertices.

5. **A strict norm improvement and a resolved extremizer question.** If

$$c_d^{\text{odd}} := \sup_{\substack{G \text{ finite abelian, } |G| \text{ odd} \\ f: G \rightarrow \mathbb{R}, \mathbb{E}f=0, f \neq 0}} \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4},$$

then

$$\frac{4804}{5465198316262956^{1/4}} \leq c_4^{\text{odd}} \leq \frac{2}{3}. \quad (\text{F.39.1})$$

The lower endpoint is approximately 0.558729391 and the upper endpoint 0.666666667. Every single inverse Fourier pair has ratio at most $2/222^{1/4} \simeq 0.518133595$, so it cannot be extremal for this problem.

6. **A classification and a sharp cost of positivity.** For an exact full-support cyclic quotient of any order m , its moment is nonnegative for every real centered function precisely when the quotient coordinates can be paired as $(a, -a)$. The minimum such positive support satisfies

$$h_{\text{pos}}(m) \sim \frac{4 \log m}{\log \log m}.$$

Thus universal positivity has asymptotically twice the support cost of an unrestricted cyclic detector.

7. **A sharp nonuniform rigidity threshold.** In the finite-fiber directional model of source 07, the optimal universal strict threshold is $[2 + (2k - 2)\kappa]^{-1}$ for Freiman order k and fiber ratio at most κ . Sharpness already holds with zero boundary loss.

Question	Audited prior result	Result supplied here
First prime torsion	$h(2) = 4, h(3) = 5, h(5) = 6;$ $h(p) \leq p + 2$	Exact determinant reduction; $p = 7, 11, 13, 17$; optimal leading asymptotic
Cyclic quotients	An isolated prime construction of linear size	Exact positivity classification; optimal cyclic support, with leading constants 2 without positivity and 4 with it
Consecutive binary determinants	Shitov’s exponent $\frac{1}{2}N \log N - 7N \log \log N$	Secondary loss $N \log \log N + O(N)$ and a cyclic Smith factor
Odd fourth-norm gap	$c_4^{\text{odd}} \leq (8/35)^{1/4}$; single-pair extremality left open	$c_4^{\text{odd}} \leq 2/3$; exact counterexample to single-pair extremality
Gowers’s cube lemma	Exact quartic coefficient in earlier reports	Smaller fourth-norm bound on that coefficient in dimension four

Table F.39.1: The comparison is with identified sources, including the incoming moment report. Earlier methods and results retain their attribution.

F.39.2 Prior results and the meaning of novelty

The determinant method is not presented as a new invention: Alon–Vu provide the large inverse entry in a unimodular binary matrix [144], and Shitov supplies the composition strategy [147]. The moment method also precedes this report [156]. The contributions are the exact transfer, the retained cyclic invariant and quantitative refinement, the explicit certificates, and the new dual-function consequences. Targeted primary-source searches did not locate these conclusions in the same form. That search is not a certificate of historical priority. Theorems are supported by their proofs, independently of any priority claim.

The results refine local cube and rigidity estimates. A new bound for the global extremal quantity $r_k(N)$ would require propagating the changed estimates through a complete density-increment iteration. That calculation remains a separate research task.

[write] **6 October 2026; sources 34–39.** *Identifications and credit within this report.* [155] is this report at its second write (sources 01–10; `article.tex` at the pin 12da35673): the sharp U^3/U^2 comparison is source 04’s Theorem III.3.4, the nonuniform-threshold question is source 07’s Question 1 (item II.9) and the prime-onset questions are source 09’s Research question on the prime-onset function (item III.16). [156] is source 29 and [157] source 31. *Not cited, and credited here:* source 28, whose archive was in `docs/incoming/` at the pin, already proves the two-harmonic lower example on $\mathbb{Z}/29\mathbb{Z}$ with the same polynomial and witness (its Corollary III.70.8), the dual identity $D_4\phi_1 = 111\phi_1 + 27\phi_3 + \phi_5$ with $\|\phi_1\|_{U^4}^{16} = 222$ (its Lemma III.70.3) and that a single inverse pair is not extremal (its Theorem III.70.9). So the table row “single-pair extremality left open” and the claim, in Section III.71, that the theorem “improves both endpoints” are accurate only against source 29’s bracket. What is new in source 34’s Section 10 is the upper bound $c_4^{\text{odd}} \leq 2/3$, the coupling of the moment profile with the dual certificate, and the refinement for $d = 5, \dots, 8$ (notes there). Checked at intake: the matrices, determinants and certificates of Section 4 and the rational comparisons of Section 10 (exact).

F.40 Source 35. The research target and the principal conclusions

Source 35, delivered Section 1.

Added in the seventh write, 6 October 2026. Delivered title page. “Additive combinatorics / Research manuscript / Graded Fourier Selection and Two Further Refinements of Gowers’ Argument / Reduced field-size thresholds, a stronger real U^4 gap, and exact polynomial partition obstructions / Prepared for Vladimir Reshetnikov and the ProveIt repository / Mathematical development and computational checks with ChatGPT / 6 October 2026”; *the PDF author field reads “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov”*. Archive `Gowers_Graded_Refinements_TeX_and_PDF.zip` (arrival 70c6a1cf1), manuscript `gowers_graded_refinements.tex` (3,007 lines, 40 pages), pin 12da35673 (its `SOURCE_AUDIT.txt`, shipped as `35-graded-selection-SOURCE_AUDIT.txt`). Printed below: the abstract, the closing lines of its title page and Section 1; its four delivered `\part` headings are printed, unnumbered, at the head of the sections they introduce: Part 1 (Sections 2–5) is a chapter of Part II (Section II.144), Part 2 (Section 6) continues the chapter on the odd-order constants in Part III (Section III.72), Part 3 (Section 7) is a chapter of Part I (Section I.60), Part 4 (Sections 8–9) is in Part V. Its macros `\pin` and `\shortpin` are renamed `\pinthirtyfive` and `\shortpinthirtyfive`; its ‘question’ environment (“Research question”, on the shared counter) is set in this report’s ‘Research question’ style.

Abstract of source 35 (as delivered). We prove three quantitative refinements related to Gowers’ proof of Szemerédi’s theorem, with the main emphasis on the arrangement-selection step. For d -arrangements of k -cubes, write $m = 2d \geq 4$ and $s = m2^k$. Weighting Fourier certificates according to their actual probability, and imposing their necessary degree balances before estimating their mass, gives the sufficient field-size exponent

$$E_{k,m} = \max_{0 \leq t \leq k} \frac{m \sum_{i=0}^t \binom{k}{i} - \binom{k}{t}}{t+1}.$$

The retained-arrangement guarantee has powers $\alpha^s \eta^{s-1}$, as in the recent one-stage Fejér refinement, while its convenient field-size exponent $2s$ is replaced by $E_{k,m} \sim 2s/k$. For eight-arrangements and $k = 2$, the comparison is 128 versus 23. The proof handles repeated physical vertices and works over finite extension fields of sufficient characteristic.

A second argument retains additional positive Fourier energy configurations and proves $\|f\|_{U^4}^{16} \geq 5 \|f\|_{U^2}^{16}$ for real mean-zero functions on finite abelian groups of odd order, improving the strongest inspected U^4 coefficient 35/8. A third argument constructs polynomial phases for which every sufficiently long flat progression has step divisible by an explicit integer D . For an allowed length set with numerical-semigroup conductor C , the exact eventual full-partition threshold is DC . This answers the inspected report’s higher-degree divisibility question and strengthens the corresponding necessary exponent for prescribed lengths. Full proofs, exact finite checks, limitations, and a research agenda are included. These results improve specified local statements; they do not supply a new global bound for progression-free sets.

Source snapshot: ProveIt, commit 12da356730. Conventional proofs with executable exact checks. No Lean verification, external refereeing, or established priority is claimed.

Gowers’ proof [1] turns a failure of uniformity into progressively more coherent algebraic structure and then into a density increment. Its quantitative cost is distributed among several distinct steps. A useful refinement must identify which step changes, preserve the normalization needed by the next one, and state the new largeness conditions explicitly. This article carries out that program for arrangement selection and gives two additional, independently proved results on the analytic and partition components.

The comparison is with the supplied ProveIt transcription and research collection, fixed at commit 12da356730b12cb73c276a258aaadf3223e3662b, together with earlier source archives whose contents were already placed in that snapshot. In particular, the comparison includes the current incoming Fejér, moment, and higher-uniformity reports, not only the merged manuscript [158, 159,

160, 163, 162]. This matters: several seemingly natural improvements, including the secondary three-cube constant, had already been obtained in those reports.

[write] **6 October 2026; sources 34–39.** *Identifications within this report.* [159] is this report at its second write (sources 01–10). [160] is source 40, `fejer-arrangement-selection.zip` (arrival 13f105f42), placed in this directory by 2dfe46f44 and not yet written in: a forward reference; its one-stage Fejér selection and the retained powers $\alpha^s \eta^{s-1}$ are source 40's (note after Theorem II.145.2). [161] is source 22, [162] source 28, [163] source 29, [164] source 18 and [165] source 23. Source 34, written the same day from the same pin, proves a better U^4 coefficient (note after Theorem III.72.1); neither cites the other.

F.40.1 The main advance: the mass of a relation stratum

An arrangement consists of $m = 2d$ axis-parallel k -cubes with common sides, lying on levels that satisfy one balanced relation. There are $s = m2^k$ labelled vertices. A map respects it when the corresponding alternating sum of its values vanishes. The selection problem is to retain many arrangements while making almost all retained arrangements respected.

The original repeated cosine selection has a very large parameter exponent. A recent repository report already replaces it by one Fejér selection [160]. Its simple sufficient field-size condition still pays for all short coefficient relations. We instead make two distinctions. First, a coefficient array is sampled with a product probability law whose largest atom is $1/L$, where L is the Fourier bandwidth. Second, after removing Boolean parity, the array has a definite polynomial degree t . A successful array of that degree belongs to a linear subspace of dimension

$$D_t = m \sum_{i=0}^t \binom{k}{i} - \binom{k}{t},$$

and each fixed array imposes at least $t + 1$ independent restrictions in the relevant probability estimate, apart from a precisely controlled degree-one diagonal. Combining these two facts yields a contribution of order

$$L^{D_t} q^{-t-1}$$

to the normalized error budget. The maximum of $D_t/(t + 1)$, rather than the number of coefficient arrays, therefore governs the sufficient exponent.

Theorem F.40.1 (Main selection consequence, informal notation). *Fix $k \geq 1$, $d \geq 2$, $m = 2d$, and a density $0 < \beta \leq 1$. Suppose a map on a density- β subset of $\mathbb{F}_q^{k+1}$ respects at least a proportion α of the natural arrangement upper bound. Under the explicit field-size and characteristic conditions of Corollary II.145.3, there is a subset with purity at least $1/(1 + \eta)$ and at least*

$$\frac{\alpha^s \eta^{s-1}}{2(s+2)^s} \beta^{m-1} q^{(m+1)k+m-1}$$

respected arrangements. For fixed k, m, β , the field-size condition is

$$q \geq C_{k,m,\beta} (\alpha \eta)^{-E_{k,m}}, \quad E_{k,m} = \max_{0 \leq t \leq k} \frac{D_t}{t+1},$$

together with a characteristic condition of order $(\alpha \eta)^{-1}$. For prime fields the displayed sufficient size condition below already enforces the characteristic condition.

The exact theorem, including a bandwidth that may be chosen freely, is Theorem II.145.2. The constant $C_{k,m,\beta}$ in the headline is not suppressed in the actual result.

F.40.2 Two additional results

The second result strengthens a real, centered norm comparison:

$$\|f\|_{U^4}^{16} \geq 5 \|f\|_{U^2}^{16} \quad \text{for finite odd-order groups and } \mathbb{E}f = 0.$$

The new ingredient is a positive contribution from repeated-entry additive configurations in the Fourier energy. It penalizes the third harmonic that can otherwise reduce the fourth-order norm of a cosine. The strongest inspected previous coefficient for this comparison is $35/8$, from [163]. We distinguish this improvement at U^4 from higher-order bounds already available in that report.

The third result concerns simultaneous polynomial phases. Assigning a different prime denominator to each degree makes flatness force divisibility of the step, one degree at a time. If the phases have degrees $d_1, \dots, d_\nu$, write

$$W = \sum_{i=1}^{\nu} \frac{d_i(d_i + 1)}{2}.$$

For minimum cell length L and oscillation tolerances ε_i , the construction has

$$D \asymp_{d_1, \dots, d_\nu} L^W \prod_{i=1}^{\nu} \varepsilon_i^{-d_i}.$$

Every good progression of length at least L has step divisible by D , regardless of where it starts. This converts the numerical-semigroup conductor of the permitted cell lengths into an exact partition threshold. It extends the linear construction of [165] and strengthens the polynomial endpoint obstruction of [164].

F.40.3 What these results settle

Table F.40.1 separates the proved improvements from questions that remain. An addition is described as new only relative to the inspected repository sources; the limited literature comparison does not certify first discovery.

Topic	Established here	Still open
Arrangement selection	An explicit sufficient exponent $E_{k,m}$, with full degree and collision accounting	The intrinsic optimal field-size threshold and propagation through the entire proof
Real norm comparison	The coefficient 5 in the U^4 -to- U^2 sixteenth-power inequality	The best coefficient and a classification of extremizing sequences
Polynomial phase partitions	Exact divisibility and conductor thresholds for the constructed families; a stronger necessary exponent	A matching general upper theorem in degree at least two

Table F.40.1: The precise scope of the three contributions.

For context, Leng, Sah, and Sawhney prove $r_\ell(N) \ll N \exp(-(\log \log N)^{c_\ell})$ for $\ell \geq 5$ [2]. No comparison of local exponents in this article asserts an improvement of that global theorem. Our assertions are the displayed finite statements and their proved consequences.

F.41 Source 36. The problem and the main result

Source 36, delivered Section 1.

Added in the seventh write, 6 October 2026. Delivered title block. “Sharp Symmetry Certification for Gowers Derivative Spectra / Alternating-rank bounds, lossless polarization, and robust finite-field inverse steps / Research manuscript prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt project / 6 October 2026”; *the PDF author field reads “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov”*. Archive `Gowers_Sharp_Symmetry_Certification.zip` (arrival 70c6a1cf1), manuscript `sharp_symmetry.tex` (1,555 lines, 19 pages), pin 6c2e2172b (its `SOURCE_AUDIT.md`, shipped as 36-symmetry-certification-SOURCE_AUDIT.md). Printed below: the abstract, the repository reference after it and Section 1; Sections 2–10 form a chapter of Part IV (Section IV.15), Sections 11–14 are in Part V. Its macros `\Q` (calligraphic $\mathcal{Q}$) and `\R` (the radical, an operator) are renamed `\Qthirtysix` and `\Rthirtysix`, and its `\pin` and `\shortpin` `\pinthirtysix` and `\shortpinthirtysix`; its ‘question’ environment (“Research question”, numbered on its own counter within sections) is set in this report’s ‘Research question’ style; its Research questions 13.1–13.12 keep their indices.

Abstract of source 36 (as delivered). Let $V = \mathbb{F}_p^n$, with p odd, and let a multilinear frequency selector $\sigma : V^k \rightarrow V^*$ be tested against the k -fold derivatives of a 1-bounded function f . We prove the dimension-independent sharp implication

$$\mathbb{E}_{\mathbf{h}} \left| \widehat{\partial_{\mathbf{h}} f}(\sigma(\mathbf{h})) \right|^2 > 1 - (1 - 1/p)^k \implies (\mathbf{h}, y) \mapsto \sigma(\mathbf{h})(y) \text{ is fully symmetric.}$$

The threshold is attained by nonsymmetric selectors in every dimension at least two, even when the selector is symmetric in its first k variables. A stronger estimate averages p^{-r} over the half-ranks of the alternating bilinear slices and retains the amplitudes of the frozen derivatives. The proof uses an explicit finite Heisenberg orthogonality inequality, proved here by a short eigenspace-multiplicity argument; its equality and stability structure are also derived. If $p > k+1$, symmetry yields a canonical polynomial $P(x) = \sigma(x, \dots, x)(x)/(k+1)!$ and an exact identity between the tested energy and $\|f e_p(-P)\|_{U_{k+1}}^{2^{k+1}}$, with a pointwise identity before averaging. This refines the global phase-extraction step around Proposition 17.2 of Gowers’s proof and extends it to vector spaces. We give noisy-selector, dense-domain, coset, multiaffine, and prime-power variants, together with explicit limitations in small characteristic. The proposed contribution is a sharp local criterion, not a new global Szemerédi bound or a general quantitative inverse theorem. The proofs are self-contained; computational checks are supplementary. Priority has not been established, and the manuscript is not proof-assistant verified.

Repository reference. ProveIt, Ramsey subtree, snapshot 6c2e2172bab8. All new statements are research statements, not additions to the repository’s checked theorem count. The accompanying source audit records full identifiers, literature versions, and integration instructions.

F.41.1 A precise extremal problem

Gowers’s global phase-extraction argument in Section 17 of [1] starts with large Fourier coefficients of multiplicative derivatives along a multilinear frequency pattern. Over a one-dimensional prime field the relevant multilinear form is automatically symmetric. Over a vector space that property is no longer automatic. This distinction suggests an extremal question that can be answered exactly:

How much derivative-spectrum energy can a genuinely nonsymmetric multilinear selector carry, uniformly in the dimension and in the bounded function being differentiated?

Here and throughout, “nonsymmetric” refers to all variables, including the variable on which the selected frequency is evaluated. Symmetry merely among the derivative directions is weaker.

Fix an odd prime p , an integer $k \geq 1$, and $V = \mathbb{F}_p^n$. The dual V^* consists of $\mathbb{F}_p$ -linear functionals $V \rightarrow \mathbb{F}_p$. All averages are normalized. Set $e_p(t) = \exp(2\pi it/p)$ and use the conventions

$$\partial_h f(x) = f(x+h)\overline{f(x)}, \quad (\text{F.41.1})$$

$$\widehat{g}(\xi) = \mathbb{E}_{x \in V} g(x) e_p(-\xi(x)). \quad (\text{F.41.2})$$

Write $\partial_{\mathbf{h}} = \partial_{h_1} \cdots \partial_{h_k}$. The derivatives commute. A function is 1-bounded when $|f(x)| \leq 1$ for every x .

For a multilinear $\sigma : V^k \rightarrow V^*$ define

$$\mathcal{A}_f(\sigma) = \mathbb{E}_{\mathbf{h} \in V^k} \left| \widehat{\partial_{\mathbf{h}} f}(\sigma(\mathbf{h})) \right|^2, \quad T(h_1, \dots, h_k, y) = \sigma(h_1, \dots, h_k)(y). \quad (\text{F.41.3})$$

Thus T is a scalar $(k+1)$ -linear form, and $0 \leq \mathcal{A}_f(\sigma) \leq 1$ for 1-bounded f . Put

$$\vartheta_{k,p} = 1 - (1 - 1/p)^k. \quad (\text{F.41.4})$$

Theorem F.41.1 (Sharp symmetry certification). *Let $f : V \rightarrow \mathbb{C}$ be 1-bounded and let $\sigma : V^k \rightarrow V^*$ be multilinear. If T is not fully symmetric, then*

$$\mathcal{A}_f(\sigma) \leq \vartheta_{k,p}. \quad (\text{F.41.5})$$

Consequently $\mathcal{A}_f(\sigma) > \vartheta_{k,p}$ forces full symmetry. For every odd prime p , every $k \geq 1$, and every $n \geq 2$, equality in (F.41.5) is attained by a nonsymmetric T and a unit-modulus f . The selector in this example is symmetric in its first k variables.

[write] **7 October 2026; sources 40–46.** *Characteristic two and general groups* (sources 43 and 46, Section IV.25). Source 43’s Theorem IV.36.4 proves the threshold $1 - (1 - 1/p_*(G))^k$ on every noncyclic finite abelian group, $p_*(G)$ the least p with a quotient $\mathbb{F}_p^2$, characteristic two included; for symmetric trilinear forms over $\mathbb{F}_2^n$ sources 43 and 46 determine the exact energy $2^{-r/2}$ and the integrable codimension $r/2$.

This theorem determines the supremum in the stated extremal problem; it is not a claim to have settled a previously named conjecture. Its proof will follow from a more informative rank-profile estimate. The basic constant is useful in a large-characteristic regime:

$$\vartheta_{k,p} \leq \frac{k}{p}, \quad \vartheta_{k,p} = \frac{k}{p} - \binom{k}{2} \frac{1}{p^2} + O_k(p^{-3}). \quad (\text{F.41.6})$$

In particular a fixed energy lower bound $a > 0$ certifies symmetry whenever $p > k/a$. This restriction matters: for fixed p and very small energy the criterion may say nothing.

Theorem F.41.2 (Canonical, lossless phase extraction). *Suppose that $p > k+1$ and that T is fully symmetric. Define*

$$P(x) = \frac{T(x, \dots, x)}{(k+1)!}, \quad g(x) = f(x) e_p(-P(x)). \quad (\text{F.41.7})$$

For every $\mathbf{h} \in V^k$, not just on average,

$$\left| \widehat{\partial_{\mathbf{h}} f}(\sigma(\mathbf{h})) \right| = |\mathbb{E}_x \partial_{\mathbf{h}} g(x)|. \quad (\text{F.41.8})$$

Therefore

$$\mathcal{A}_f(\sigma) = \|f e_p(-P)\|_{U_{k+1}}^{2^{k+1}}. \quad (\text{F.41.9})$$

In particular the conclusion holds if the energy exceeds $\vartheta_{k,p}$. For the identity itself neither an energy lower bound nor boundedness of f is needed.

The diagonal polynomial in this statement is elementary polarization, not a new integration mechanism. The refinement is the explicit, pointwise, equality-form interface together with a sharp test that certifies its symmetry hypothesis in arbitrary dimension.

F.41.2 What is, and is not, improved

Proposition 17.2 in the corrected ProveIt transcription assumes $k + 1 < N$ and concludes a polynomial of degree at most $k + 1$ [166]. We retain that indexing. In its global one-dimensional setting, Theorem F.41.2 gives a particular polynomial and equality, rather than selecting a phase through a mixed cube inequality. Its scalar lower bound in the parameter α already has coefficient one; we do not claim to improve that coefficient. The genuinely additional quantitative statement is Theorem F.41.1 and its rank-profile refinement over vector spaces.

The finite Heisenberg representation theory behind the proof is standard; see, for example, Gurevich–Hadani [149]. The present manuscript proves the needed finite-dimensional facts directly. Existing inverse theorems also address much weaker correlations and much more general intermediate forms. In particular, Tidor [150] separates symmetrization from integration in low characteristic. Our exact criterion is for the derivative array of one function and does not replace general mixed-function symmetrization.

The contemporary benchmark also includes Milićević’s September 2026 preprint [32], whose stated high-characteristic inverse theorem has quasipolynomial bounds. We do not derive a competing inverse theorem here. Nor do we change an estimate for progression-free subsets of $[N]$; for comparison, Leng–Sah–Sawhney [2] prove, for $\ell \geq 5$, bounds of the form $r_\ell(N) \ll N \exp(-(\log \log N)^{c_\ell})$. These comparisons delimit the claim, rather than supply ingredients for our proofs.

[write] **6 October 2026; sources 34–39.** *Credit within this report* (checked at intake and at the write). The bilinear inequality of Section IV.17 sharpens source 31’s Lemma II.106.4 (the noncommuting-subgroup gap, bound $1/r \leq 1/2$ with r the commutator order; proved by the same twirl and eigenspace-multiplicity argument), which source 36 does not cite although source 31’s archive was in `docs/incoming/` at its pin; the bounds coincide only for alternating rank two (half-rank one), where both are $1/p$ and $1/p$ is sharp. At $k = 1$ Theorem F.41.1 is the vector-space, sharp-in- p analogue of source 29’s symmetry step (Section II.110, “Forcing symmetry without dividing by two”, a gap $1/2$ on every finite abelian group), also not cited. Source 38 reproves the same gap for odd order (its Lemma II.150.3). The formal `proposition_17_2` (`Sections17_18.lean`, line 48) has the exact companion `proposition_17_2_holds` (`Proofs17PhaseRemoval.lean`, line 373, at HEAD 89c1cbc38), on $\mathbb{Z}/N\mathbb{Z}$; Theorem F.41.2 is an equality-form, vector-space refinement of its global scalar step and does not change its status. Sources 43 and 46, placed in this directory by 2dfe46f44 and not yet written in, treat characteristic two.

F.42 Source 37. Research target and relation to the repository

Source 37, delivered Section 1.

Added in the seventh write, 6 October 2026. Delivered title page. “Sharp Spectral Detection of Missing Intervals / Exact moment tradeoffs, extremal rigidity, and quantitative recurrence in Gowers’s argument / Research manuscript prepared with ChatGPT for Vladimir Reshetnikov / 6 October 2026”; *the PDF author field reads the same*. Archive `Gowers_Sharp_Missing_Intervals.zip` (arrival `a33bcaedf`), manuscript `sharp_missing_intervals.tex` (1,616 lines, 27 pages), pin `869289aa3` (its `SOURCE_AUDIT.md`, shipped as `37-missing-intervals-SOURCE_AUDIT.md`). Printed below: the abstract, the status statement of its title page and Section 1; Sections 2–9 and Appendices A–B form a chapter of Part I (Section I.49), Sections 10–12 and Appendix C are in Part V. Its macros `\T` (the circle $\mathbb{R}/2\pi\mathbb{Z}$), `\Prob` (calligraphic $\mathcal{P}$, a set of measures), `\code` (`\nolinkurl`) and `\pin` are renamed `\Tthirtyseven`, `\Probthirtyseven`, `\codethirtyseven` and `\pinthirtyseven`; its ‘question’ environment (“Question”, numbered on its own counter within sections) is set in this report’s ‘Question’ style, and its Questions 11.1–11.10 keep their indices. Its definitions are set upright, in this report’s style (it sets them in italics).

Abstract of source 37 (as delivered). A missing interval forces a low Fourier coefficient. We solve an exact version of this problem in the first nontrivial spectral window, including the complete tradeoff between interval mass and a uniform bound on finitely many Fourier moments. For a probability measure on the circle and $\pi/(R+1) < a < 2\pi/(R+1)$, the optimal lower bound for the largest of its first R Fourier coefficients, when the open arc $(-a, a)$ is empty, is

$$\rho_R(a) = \frac{1}{R - \cot(a/2) \tan((R+1)a/2)}.$$

The proof gives a positive-coefficient polynomial certificate, a unique $R+1$ -point extremizer, exact noisy-gap extremizers, and stability estimates. A finite-cyclic consequence replaces the frequency range $(N/M)^2$ in the proof-supported version of Gowers’s Lemma 5.2 by $\lfloor N/M \rfloor - 1$, and replaces the coefficient bound $tM/(4N)$ by $tM/(N-M)$. Using the same Weyl input, a separate optimization of the Dirichlet denominator improves the monomial-recurrence exponent in Lemma 5.5. All new assertions are proved in the manuscript; the quantitative Weyl estimate is an explicitly identified input. Numerical and symbolic regression checks, a repository integration plan, and further research problems accompany the proofs.

Status and scope. This is an unrefereed mathematical research draft, not a Lean-verified result. The exact formulas and refinements are developed and proved here, but their historical priority is not established. Classical positive-polynomial and truncated-moment machinery is credited explicitly. No improved global Szemerédi bound or record Diophantine-approximation bound is claimed.

Intended integration: `Combinatorics/Ramsey/Research/GowersSzemeredi/`

F.42.1 A local bottleneck with an exact extremal formulation

The target is the missing-interval step in Section 5 of Gowers’s proof [1]. A positive measure supported away from zero must have a nontrivial low Fourier moment. Gowers uses an interval convolution and a Fourier tail estimate. The present approach constructs a genuinely bandlimited polynomial with the required sign pattern. The resulting improvement is not just a change in a tail constant: the frequency scale changes from quadratic to linear in the inverse gap radius.

The natural optimization problem is more informative than the single lemma. Given a gap radius a and a cutoff R , how small can all of the first R Fourier moments be? If those moments are merely small rather than forced by an empty gap, how much mass must lie in the gap? We solve both questions exactly when $\pi < (R+1)a < 2\pi$. We also identify the equality measures, prove a quantitative rigidity statement, and derive a recurrence consequence using an independently adjustable Dirichlet scale.

F.42.2 The precise baseline

The repository version of Section 5 distinguishes a printed assertion from what its proof supports [167]. Its formal target for Lemma 5.2 uses

$$0 < |r|_N \leq \frac{N^2}{M^2}, \quad |\widehat{1_A}(r)| \geq \frac{|A|M}{4N}, \quad (\text{F.42.1})$$

under the assumptions $0 < M$, M even, $2M \leq N$, and $A \cap [-M, M) = \emptyset$ in $\mathbb{Z}/N\mathbb{Z}$. Here the Fourier transform is unnormalized and $|r|_N$ is the least absolute residue. The printed statement has $1/2$ where the displayed proof gives $1/4$. We compare with the proof-supported $1/4$ version. This does *not* assert that the printed $1/2$ inequality is false: the stronger theorem below implies it.

The source file `Proofs05PolynomialPartition.lean` already handles colliding polynomial values with a multiplicity function. We retain that correction and do not present it as a new contribution. Likewise, the repository’s squared threshold for square-root-length recurrence is respected; we do not substitute an unrounded real square root into a theorem whose hypothesis is an integer sample length.

The main inspected Section 5 files are pinned at

869289aa38d59354e8dac2213df08ae0dd256415.

The research-directory README was also inspected for overlap. It describes many earlier refinements of density extraction, cube counts, arrangement selection, and related inverse steps. The present manuscript instead focuses on the missing-interval extremal problem and its recurrence consequence. That comparison is an overlap audit of the inspected material, not an assertion that all unpublished work or all classical literature has been exhausted.

F.42.3 What is proved, and what is imported

The main exact theorem is Theorem I.50.1. Its lower bound uses an explicit trigonometric polynomial. Its attainment uses a finite-dimensional moment criterion proved in Section I.52; this is classical positive-polynomial territory, closely related to the frequency-selective Toeplitz decomposition of Yang and Xie [152]. The explicit extremal formula, equality analysis, and application are the focus here, not a claim to have invented the general moment criterion.

The recurrence argument imports the Weyl inequality in a precisely stated form. It does not claim to improve that inequality. Work such as Baker’s on small fractional parts [153] and Maynard’s on simultaneous polynomial recurrence [72] belongs to a stronger and broader Diophantine-approximation literature. Our recurrence theorem is a transparent improvement of the selected local argument, not a claim of a best-known exponent. Propagating it through the polynomial partition and the later Szemerédi induction is a further problem, not a completed consequence.

[write] **6 October 2026; sources 34–39.** *Checked at HEAD 89c1cbc38.* The baseline (F.42.1) is the formal `lemma_5_2` (`Section05.lean`, line 141): unnormalized `fourier (ZMod.dft)`, $0 < M$, M even, $2M \leq N$, A disjoint from `centeredIntervalNM` $[-M, M)$, conclusion $0 < |r|_N \leq N^2/M^2$ and $tM/(4N) \leq |\widehat{1_A}(r)|$; its exact companion is `lemma_5_2_holds` (`Proofs05MissingInterval.lean`, line 220). `lemma_5_3` (line 152), `lemma_5_5` (line 181) and `lemma_5_5_square_root_auxiliary` (line 192) are as described, with exact companions `lemma_5_3_holds` (`Proofs05Weyl.lean`, line 3607) and `lemma_5_5_holds` (`Proofs05PolynomialPartition.lean`, line 1678). The three Lean files the source cites have the blobs it records. No earlier source of this report treats Lemma 5.2; this is the first.

F.43 Source 38. Research question, main results, and provenance

Source 38, delivered Section 1.

Added in the seventh write, 6 October 2026. Delivered title block. “Exact Quotient Reduction and Sharp Stability in Gowers’ Ramsey Framework / Characteristic-five cube extremizers, odd-order rigidity, and six-term norm-extension defects / Prepared for Vladimir Reshetnikov / Mathematical development and verification with ChatGPT / 6 October 2026”; *the PDF author field reads “Mathematical research developed with ChatGPT for Vladimir Reshetnikov”*. Archive `gowers_quotients_rigidity.zip` (arrival `b103a65d4`), manuscript `gowers_quotients_rigidity.tex` (3,217 lines, 46 pages), pin `12da35673` (its `SOURCE_LEDGER.json`, shipped as `data/38-quotients-SOURCE_LEDGER.json`). Printed below: the abstract and Section 1; its four delivered `\part` headings are printed, unnumbered, at the head of the sections they introduce: Part I (Sections 2–5) is a chapter of Part III (Section III.82), Part II (Section 6) a chapter of Part II (Section II.149), Part III (Section 7) a chapter of Part III (Section III.87), Part IV (Sections 8–10) is in Part V. Its macro `\Q` (the cube count, a calligraphic $\mathcal{C}_3$) is renamed `\Qthirtyeight`, and its `\ip` (no `\left/\right)` is set as source 31’s `\ipthirtyone`, which has the same definition; its figure path is renamed to the shipped file.

Abstract of source 38 (as delivered). We resolve three local questions left open by the preceding Ramsey research reports. First, on every nontrivial finite abelian group of exponent five, the maximum three-cube count of a real perturbation of a positive constant, at fixed centered U^3 norm u , reduces exactly to a five-point quotient whenever $u/\delta \leq 10^{-5}$. Every maximizer has this quotient form. The proof retains a coercive Fourier energy defect and shows that the sixth- and seventh-degree errors are linear in the norm lost under projection. This settles the formerly undetermined dimension-uniform eighth-order coefficient at $20/9$. We additionally obtain the tenth- and twelfth-order coefficients $869/(216 \cdot 2^{3/4})$ and $65759/(5832 \cdot 2^{3/2})$, an even analytic extremal profile, and quantitative control of the off-quotient remainder. More generally, for every fixed odd exponent m , exact reduction to a cyclic quotient holds when $u/\delta \leq 10^{-8}m^{-8}$.

Second, we determine the sharp odd-order local stability modulus for the critical U^3 – L^2 inequality through its first fractional term: $\varepsilon/4 + 3\varepsilon^2/16 + (\sqrt{2}/64)\varepsilon^{5/2} + O(\varepsilon^3)$. The fractional term is forced by order-three characters and disappears when the group order is coprime to six. A finite Weyl-coset expansion gives an explicit uniform error bound.

Third, for a layered linear–cubic–quintic norm coloring plus an arbitrary function of independent coordinates, we calculate the exact number of symmetric six-term progressions. Its excess above the basic surjectivity count is governed by centered third and fifth differences and a quadratic-character statistic. All three results have complete proofs, sharpness or equality analysis, and reproducible checks. They refine local parts of Gowers’ framework; they do not produce a new global bound in Szemerédi’s theorem.

The supplied Ramsey project contains both a reconstruction of Gowers’ proof of Szemerédi’s theorem and a growing collection of local refinements [1, 168, 169]. A useful next contribution must do more than repeat the existing density-selection estimates or improve a constant already computed in another report. We therefore start from three explicitly unresolved questions: whether characteristic-five cube extremizers depend only on a five-point quotient, what the optimal odd-order stability coefficients are, and how higher norm layers interact with independent additions to a coloring.

The first question has the strongest answer here. The ambient number of variables disappears from the entire small-norm optimization problem. The equality of maxima is exact, has a dimension-independent numerical range, and includes a classification of all maximizing functions by a quotient. The second result identifies a new fractional-power correction rather than only improving a rough integer-power bound. The third supplies an exact formula for the first extension problem where two nonlinear norm layers interact.

F.43.1 Three distinct extremal problems

For a finite abelian group G , all spatial averages are normalized. For real W put

$$\mathcal{C}_3(W) = \mathbb{E}_{x, h_1, h_2, h_3 \in G} \prod_{\omega \in \{0,1\}^3} W(x + \omega \cdot h) = \|W\|_{U^3}^8.$$

For a nontrivial G and $\delta, u > 0$, define

$$\mathcal{M}_{\delta, G}(u) = \max\{\mathcal{C}_3(\delta + f) : f : G \rightarrow \mathbb{R}, \mathbb{E}f = 0, \|f\|_{U^3} = u\}.$$

The maximum exists because the norm sphere in this finite-dimensional real space is compact.

Theorem F.43.1 (Main quotient and expansion theorem). *If G is nontrivial and $5G = 0$, then*

$$\mathcal{M}_{\delta, G}(u) = \mathcal{M}_{\delta, \mathbb{Z}/5\mathbb{Z}}(u) \quad (0 < u \leq 10^{-5}\delta).$$

Every maximizer factors through a nonzero homomorphism $G \rightarrow \mathbb{Z}/5\mathbb{Z}$. As $u/\delta \rightarrow 0$,

$$\begin{aligned} \mathcal{M}_{\delta, G}(u) = & \delta^8 + 6\sqrt{2}\delta^4 u^4 + \frac{2^{3/4}}{3}\delta^2 u^6 + \frac{20}{9}u^8 \\ & + \frac{869}{216 \cdot 2^{3/4}} \frac{u^{10}}{\delta^2} + \frac{65759}{5832 \cdot 2^{3/2}} \frac{u^{12}}{\delta^4} + O\left(\frac{u^{14}}{\delta^6}\right). \end{aligned} \quad (\text{F.43.1})$$

The implicit constant is absolute and independent of G . More generally, the extremal value is given near zero by a convergent power series in $(u/\delta)^2$, multiplied by δ^8 . For $0 < \delta < 1$, the same maximum is attainable with $0 \leq \delta + f \leq 1$ if also $4u \leq \min\{\delta, 1 - \delta\}$.

The exact reduction is proved in Section III.84; the analytic calculation, including its phase optimization, is in Section III.85. The small numerical cutoff is an explicit sufficient certificate; its optimal value remains to be determined. The convergent series is asserted on a possibly smaller absolute neighborhood of zero, as is appropriate for the implicit function theorem.

The quotient principle extends beyond exponent five. If G has fixed odd exponent $m \geq 3$, then

$$\mathcal{M}_{\delta, G}(u) = \mathcal{M}_{\delta, \mathbb{Z}/m\mathbb{Z}}(u) \quad \text{for } 0 < u \leq 10^{-8}m^{-8}\delta.$$

Every maximizing function factors through some cyclic quotient of order dividing m . This broader result, proved after the five-point profile, separates the arithmetic of a single cyclic group from the ambient number of variables. The dependence of the sufficient cutoff on m is not claimed optimal.

A separate problem concerns large uniformity norm relative to L^2 . Let $\mathcal{Q}_{\cos}(G)$ be the normalized quadratic phases on subgroup cosets, as defined precisely in Part II. For $\|v\|_2 = 1$, set

$$\kappa(v) = \|v\|_{U^3}^8, \quad \mathcal{F}(v) = \max_{\psi \in \mathcal{Q}_{\cos}(G)} |\langle v, \psi \rangle|^2.$$

Thus $\mathcal{F}$ is squared correlation, sometimes called fidelity. Define the worst infidelity at deficit at most ε by

$$\mathfrak{S}_{\text{odd}}(\varepsilon) = \sup_{\substack{|G| \text{ odd} \\ \|v\|_2 = 1 \\ \kappa(v) \geq 1 - \varepsilon}} (1 - \mathcal{F}(v)).$$

Theorem F.43.2 (Main odd-order stability theorem). *Uniformly over finite abelian groups of odd order,*

$$\mathfrak{S}_{\text{odd}}(\varepsilon) = \frac{\varepsilon}{4} + \frac{3\varepsilon^2}{16} + \frac{\sqrt{2}}{64}\varepsilon^{5/2} + O(\varepsilon^3).$$

For the supremum restricted to $\gcd(|G|, 6) = 1$, the corresponding formula is

$$\mathfrak{S}_{6'}(\varepsilon) = \frac{\varepsilon}{4} + \frac{3\varepsilon^2}{16} + O(\varepsilon^3).$$

Both assertions are sharp through the displayed orders.

An explicit finite estimate and matching examples are proved in Section II.150. Here the parameter measures proximity to a *maximal* norm under an L^2 constraint. It is unrelated to the small centered norm in Theorem F.43.1.

The third result concerns a finite field $\mathbb{F}_q$ of characteristic greater than five. Let $N_j : \mathbb{F}_{q^j} \rightarrow \mathbb{F}_q$ be the field norm and let $R : \mathbb{F}_q^e \rightarrow \mathbb{F}_q$ be arbitrary. For

$$P(a, u, v, w) = a + N_3(u) + N_5(v) + R(w),$$

we count all nonzero directions for which the six values at a center plus $\pm h/2, \pm 3h/2, \pm 5h/2$ agree in opposite pairs. Theorem III.88.2 gives the exact count in terms of third and fifth differences of R , with a quadratic-character correction. Equality in the earlier surjectivity bound is classified by additive quadraticity of R . These are symmetric color equalities; the six colors need not all coincide.

F.43.2 What is added to the preceding reports

Question	Previously established	Addition in this article
Characteristic-five cube profile	Sharp u^6 coefficient; 20/9 at order u^8 on $\mathbb{Z}/5\mathbb{Z}$ only [170, 171]	Exact quotient reduction, all ambient dimensions, u^{10} and u^{12} coefficients, and a theorem for every fixed odd exponent
Odd-order critical-norm stability	An envelope with coefficient 12; its optimal odd-order second coefficient left open [172]	Sharp asymptotic coefficient 6, the order-three fractional correction, and a uniform finite remainder
Independent additions to norm colorings	Surjectivity for complete norm cores and an exact cubic-core defect [173]	Exact six-term count for the cubic–quintic core, including the character correction and zero-defect classification

The old five-point coefficient 20/9 is not presented as newly calculated here. The new conclusion is that it governs the unrestricted dimension-uniform problem, because all exact maximizers reduce to five points. The tenth- and twelfth-order coefficients and the exact quotient theorems are additional results. Likewise, the norm-layer construction itself is due to earlier work; our addition is its exact response to an arbitrary independent extension.

F.43.3 Literature and scope

Gowers’ Lemmas 3.8–3.10 supply cube Cauchy–Schwarz, the norm triangle inequality, and an upper cube-count bound [1]. Part I sharpens the last of these in a specified finite-field regime. The underlying norm-geometric perspective is consistent with Hatami’s work on generalizations of Gowers norms [31]. For the critical U^3 – L^2 problem, the classification of extremizers and qualitative near-extremizers is established by Eisner–Tao [7, Theorem 1.4]. Weyl operators and isotropic subgroups also occur in stabilizer testing [125]; our calculation determines a particular sharp asymptotic modulus within that framework.

For symmetric progression colorings, Shi–Dong’s layered norm method is the relevant primary source [33]. Their paper also explains why the historical conjectures in Gowers’ Section 4 cannot simply be treated as unresolved in 2026. We use the construction as prior work and study an exact extension question, without claiming a better coloring exponent.

The present results do not improve the global size bound for progression-free subsets of the integers. For orientation, the Leng–Sah–Sawhney theorem gives, for fixed $k \geq 5$,

$$r_k(N) \ll_k N \exp(-(\log \log N)^{c_k}), \quad c_k > 0 \quad [2].$$

A local sharp coefficient does not imply an improvement of this theorem or of the complete historical density-increment argument. In particular, an exponent-five group has no arithmetic progression with more than five distinct points, although its cube norms remain well defined.

The claims here are mathematical statements with proofs and finite verification, developed with AI assistance. Independent checks in this package are separate AI derivations and computations, not external refereeing or Lean kernel certification. The source search establishes the relation to the inspected reports and papers; it does not establish exhaustive bibliographic priority.

[write] **6 October 2026; sources 34–39.** *Identifications within this report.* [169] is this report at its second write (sources 01–10); [170] is source 21, [171] source 23, [172] source 31 and [173] source 22. The three “explicitly unresolved questions” are source 21’s Questions on the universal eighth-order coefficient, fixed higher dimensions and further terms on the five-point group (item III.27), source 31’s odd-order second coefficient (item II.46) and the extension of source 22’s norm-core defects (item III.28). The credits in the table above are accurate (checked at intake).

F.43.4 Source checkpoint and normalization

The repository was inspected through the GitHub connection at

12da356730b12cb73c276a258aaadf3223e3662b.

The files include the edited Gowers transcription, the consolidated research article, the formalization status, and the Section 3 declarations and companion proofs. Several later reports were retrieved separately because their mathematics had not yet been incorporated into that consolidated article. Their names and content hashes are recorded in the accompanying source ledger.

Part	Normalization	Quantity optimized
I	$\mathbb{E}f = 0, \ f\ _{U^3} = u$	Cube count of $\delta + f$
II	$\ v\ _2 = 1, \ v\ _{U^3}^8 \geq 1 - \varepsilon$	Distance from quadratic coset phases
III	Counting measure on centers and directions	Symmetric six-term progressions

Zero directions and repeated vertices are included in all cube averages. In Part III zero directions are explicitly subtracted from the progression count. All dual sums use counting measure, in contrast to the normalized spatial averages.

F.44 Source 39. The local problem and the scope of the improvement

Source 39, delivered Section 1.

Added in the seventh write, 6 October 2026. Delivered title block. “Optimal Shared-Anchor Interpolation and Labelled Slice Refinement in Gowers’s Szemerédi Argument / Exact finite minimax bounds, sharp Poisson limits, and a replacement for the sampling step of Lemma 16.10 / Research manuscript prepared with ChatGPT for Vladimir Reshetnikov / 6 October 2026”; *the PDF author field reads the same*. Archive `ramsey_shared_anchor_interpolation.zip` (arrival `b103a65d4`), manuscript `article.tex` (2,164 lines, 31 pages), pin `869289aa3` (its `SOURCE_AUDIT.md`, shipped as `39-shared-anchor-SOURCE_AUDIT.md`). Printed below: the abstract, its keywords and Section 1; Sections 2–11 form a chapter of Part II (Section II.151), Sections 12–15 and Appendices A–B are in Part V. Its macro `\file` (typewriter type, detokenized) is set as this report’s `\filename`, which has the same definition, and its `\pin` is renamed `\pinthirtynine`; its ‘question’ environment (“Research question”, on the shared counter) is set in this report’s ‘Research question’ style.

Abstract of source 39 (as delivered). A local step in Gowers’s proof of Szemerédi’s theorem selects common column coordinates from which affine functions on many different fibres can be reconstructed. We replace size truncation followed by repeated sampling with uniform sampling of a fixed-size subset, retaining targets that are themselves sampled. For fibres partitioned into at most q reconstruction classes and a threshold of t distinct anchors, we prove that r shared anchors leave at most $tq(m - r)/(m(r + 1))$ of the ambient mass uncertified, where m is the number of columns. In particular, affine reconstruction has an $O(q/\sigma)$ anchor budget at exceptional mass σ , without a minimum class size or a separate long-column hypothesis.

We determine the exact finite minimax loss, prove that balanced class sizes capped at the first one-class maximum attain it, and obtain an explicit quadratic formula for the affine cap. Matching permutation-orbit examples and an exact conditional-expectation algorithm are supplied. A uniform hypergeometric-to-Poisson argument yields the complete asymptotic profile for $r \rightarrow \infty$ and $r/m \rightarrow 0$. For affine reconstruction in the low-loss regime, the sharp coefficient is $\varphi^3 e^{-\varphi} = 0.8399620946\dots$, where φ is the golden ratio. We also prove higher-degree and interpolation-space extensions, an additive corruption bound, and a labelled slice-refinement theorem that avoids replacing each individual slice list by the entire union list.

The results are local to the interpolation mechanism. Their finite minimax optimality concerns the explicitly defined certification rule, not unrestricted algebraic reconstruction. We compare the new bounds with the inspected **ProveIt** development, including its existing sampled-or-anchored interface and geometric synchronization. The separate closing-parameter obstruction in Lemma 16.10 is not resolved. Complete mathematical proofs and reproducible exact-arithmetic checks are supplied; no new Lean theorem is claimed to have been kernel-checked, and priority has not been established.

Keywords. Szemerédi’s theorem; shared sampling; hypergeometric occupancy; minimax bounds; interpolation; multilinear graph covers.

F.44.1 Why the interpolation step is worth isolating

Gowers’s argument [1] passes through functions that are approximately polynomial in several variables. In the step surrounding Lemma 16.10, a function of a base variable h and a final coordinate x is covered, on each vertical fibre, by a bounded number of affine graphs. One seeks a small set of coordinates $a_1, \dots, a_r$ that can serve as interpolation anchors for many fibres simultaneously. The horizontal sections at those coordinates can then be refined and lifted to multilinear candidates in (h, x) .

The important word is *simultaneously*. Selecting two points separately in each fibre does not give a bounded family of horizontal sections. Conversely, requiring every fibre class to be hit by a common sample is unnecessarily strong. The argument only needs to discard a small amount of total ambient mass. That distinction makes a size-weighted sampling estimate more effective than a worst-class failure estimate followed by a union bound.

The printed argument first removes small classes, then samples with replacement. Its estimate for the chance of obtaining two distinct points uses a two-hit expression that does not account for repeated draws of the same point. The inspected **ProveIt** development already records this obstruction and proves corrected distinct-anchor sampling statements [174, 175]. The counterexample and its repair are not new results of this article.

The repository also already handles short column sets by sampling all columns and allowing a target to survive when its own column is sampled [176]. Its later files perform a common slice refinement and synchronize suitable rectangular products into proper common-step boxes [177, 178, 179]. We use these facts to identify precisely what is improved here: the *mass-weighted shared sample*, its *exact extremal theory*, and the *preservation of individual slice lists*. We do not claim that direct sampling or synchronized interpolation is new to the repository.

F.44.2 The principal statements

Write J for an m -element set of columns. In row h , let $B_h \subseteq J$ be the domain of interest, partitioned into at most q classes. A target is certified if it lies in the sample itself, or if its class contains t distinct sampled columns. For affine functions $t = 2$; for polynomials of degree at most d , take $t = d + 1$.

The basic bound, proved in Section II.153, is

$$\min_{|R|=r} \frac{1}{m|H|} \sum_{h \in H} |B_h \setminus K_h(R)| \leq \frac{tq(m-r)}{m(r+1)}. \quad (\text{F.44.1})$$

Here $K_h(R)$ is the certified set. The same statement holds for arbitrary nonnegative row weights, with q replaced by the weighted mean number of classes. The minimum is over one set R shared by every row.

The finite answer is more precise. Define

$$\ell_{m,r,t}(s) = \frac{1}{m \binom{m}{r}} \sum_{j=0}^{\min(s,r,t-1)} (s-j) \binom{s}{j} \binom{m-s}{r-j}. \quad (\text{F.44.2})$$

Then the worst possible best-anchor loss is exactly

$$\mathfrak{M}_{m,q,r,t} = \max_{\substack{s_1, \dots, s_q \in \mathbb{N}_0 \\ s_1 + \dots + s_q \leq m}} \sum_{i=1}^q \ell_{m,r,t}(s_i). \quad (\text{F.44.3})$$

The upper bound is an averaging argument. For the reverse inequality we build a family containing all coordinate permutations of an extremizing class profile; every anchor set then has the same loss. Section II.155 goes further: an optimal profile is balanced and capped at the first maximum of ℓ , and the affine cap has an explicit quadratic-radical formula. Thus the finite minimax problem is solved without an optimization oracle.

Set

$$F_t(x) = xe^{-x} \sum_{j=0}^{t-1} \frac{x^j}{j!}, \quad \lambda_t = \operatorname{argmax}_{x \geq 0} F_t(x), \quad \kappa_t = F_t(\lambda_t). \quad (\text{F.44.4})$$

For fixed t , as $r \rightarrow \infty$, $r/m \rightarrow 0$, and q/r remains bounded,

$$\mathfrak{M}_{m,q,r,t} = \frac{q}{r} \left[F_t \left(\min \left\{ \lambda_t, \frac{r}{q} \right\} \right) + o(1) \right]. \quad (\text{F.44.5})$$

In particular, if also $q/r \rightarrow 0$, then $\mathfrak{M}_{m,q,r,t} \sim \kappa_t q/r$. For $t = 2$,

$$\lambda_2 = \varphi = \frac{1 + \sqrt{5}}{2}, \quad \kappa_2 = \varphi^3 e^{-\varphi}. \quad (\text{F.44.6})$$

Thus the order q/σ of the sufficient affine anchor budget is optimal for this class-based certification mechanism in a sparse-sampling regime.

F.44.3 What is, and is not, being optimized

The minimax game fixes a recovery rule. It does not charge computation time, and it does not allow extra algebraic information to certify a point that the rule leaves uncertified. An interpolation routine might infer a polynomial from other identities, notice that two labelled classes carry the same polynomial, or exploit special geometric restrictions on the row family. Those possibilities are deliberately outside the lower-bound model.

This distinction is particularly important in applying the results to Gowers’s proof: the extremizing permutation-orbit examples are arbitrary families of class partitions. We do not show that they satisfy every other hypothesis of the full inverse argument. Our sharpness assertion is a sharp statement about a local reusable lemma, not a lower bound for Szemerédi numbers or for the entire inverse theorem.

Occupancy and missing-mass optimization are established subjects. Berend and Kontorovich study extremal expected missing mass [75]; low-frequency occupancy masses also occur in work on Good–Turing estimation [154]. We therefore make no claim to originate occupancy methods. The package here connects a finite, without-replacement, partial-domain occupancy problem to shared interpolation, proves its exact minimax realization, and tracks the resulting labelled graph-cover parameters. All arguments needed for our results are included below.

[write] 6 October 2026; sources 34–39. *Credit within this report: source 19.* Source 39 re-derives, without citing it, the sampling theory of source 19 (Sections II.67–II.69), which was placed in this repository (58aa0c493) before source 39’s pin: the common r -subset sampler with an $O(q/\sigma)$ budget (source 19’s Lemma II.67.4: $qd/(r+1)$), the exact finite minimax as a maximum over class profiles, attained by permutation orbits (Theorem II.68.2), its dynamic program, deterministic selection by conditional expectation, the sharp sparse law with the affine constant $\varphi^3 e^{-\varphi} = 0.839962094657$ (Corollary II.69.5) and the elementary-symmetric candidate count. The textual overlap is below 0.5% (intake, 8-word shingles); the overlap is of results. *The two certification rules differ.* Source 19 recovers a class C when $|R \cap C| \geq \min(d, |C|)$ and otherwise discards the whole class, sampled points included (loss $s \mathbf{1}\{a < \min(d, s)\}$ for a class of size s meeting the sample

in a points). Source 39 keeps every sampled point and discards only the unsampled points of a class with fewer than t anchors (loss $(s - a)\mathbf{1}\{a < t\}$, Definition II.152.2); here t is source 19’s d . New in source 39 are: this rule, with the bound $tq(m - r)/(m(r + 1))$; the closed solution of the finite profile problem for this rule (Section II.155); the two-regime limit $F_t(\min\{\lambda_t, r/q\})$ (source 19 assumes $q/r \rightarrow 0$); unisolvent spaces and additive corruption; and label-preserving common refinement. Source 39 is printed after source 19’s chapter, with source 19 as the spine (Section II.151).

F.44.4 Status and source discipline

The repository comparison uses the fixed snapshot 869289aa38d59354e8dac2213df08ae0dd256415. The precise source files and their blob identifiers are recorded in the accompanying `SOURCE_AUDIT.md`. The article is an unrefereed research draft. The proofs have been developed and checked here, and finite cases have been tested with exact rational arithmetic, but there is no assertion of an independent mathematical review or a new proof-assistant verification. The results should be reviewed on their stated hypotheses before integration as trusted lemmas.

F.45 Source 40. The local problem and the contribution

Source 40, delivered Section 1.

Added in the eighth write, 7 October 2026. Delivered title block. “One-stage Fejér selection for Gowers arrangements / Single-monomial relation isolation, polynomial parameter loss, and sharp exponents within a positive-Fourier selection class / Research note prepared for the ProveIt repository / October 6, 2026”; *the PDF author field reads the same; no statement on AI assistance is made.* Archive `fejér-arrangement-selection.zip` (arrival 13f105f42), manuscript `fejér-arrangement-selection.tex` (1,281 lines, 20 pages); *no pin and no blob identifiers* (its bibliography records the repository paths and “inspection date rather than ... a commit-pinned audit”). Printed below: the abstract and Section 1; Sections 2–9 form a chapter of Part II (Section II.164), Sections 10–12 are in Part V. Its ‘question’ environment (“Research question”, on the shared counter) is set in this report’s ‘Research question’ style. Source 35 builds on this source (its `PriorFejér`): that forward reference is now resolved.

Abstract of source 40 (as delivered). Section 15 of Gowers’ 2001 proof of Szemerédi’s theorem uses repeated cosine-based random selections to pass from many respected cubical arrangements to a subset in which almost every arrangement is respected. We give a self-contained alternative based on a single normalized Fejér kernel. For ordered d -arrangements of k -cubes, put $s = 2d2^k$ and $M = (2d + 1)k + 2d - 1$. If a map on a density- β subset of $\mathbb{F}_q^{k+1}$ respects at least $\alpha\beta^{2d-1}q^M$ arrangements, then, under explicit characteristic and field-size conditions, there is a subset containing at least

$$\frac{\alpha^s \eta^{s-1}}{3^{s-1}(s+1)^s} \beta^{2d-1} q^M$$

arrangements, of which at least a proportion $1 - \eta$ are respected. For $d = 8$, this replaces the exponent $s2^{s-1}$ in Gowers’ Lemma 15.5 by exponents s and $s - 1$ in the abundance and accuracy parameters. A single feature, $F(y, r) = r \prod_i y_i$, is enough: we identify its universal relation space and bound all unintended bounded-coefficient relations. The proof treats repeated vertices explicitly. Further results include vector-valued targets with dimension-free constants, an exact formula for target-torsion losses, and an exponent obstruction for homogeneous nonnegative-Fourier kernels, including independent repeated stages. These are local refinements, not a claimed new bound for Szemerédi’s theorem. All main results are proved here; finite exact checks are supplied as supplementary evidence, not as a formal verification or a claim of established priority.

F.45.1 Why focus on the selection step?

Gowers' paper [1] develops several transitions from abundant combinatorial configurations to increasingly coherent algebraic structure. This note concerns one such transition in Section 15: a map respects a positive proportion of certain cubical arrangements, and one wants to find a subset where the proportion is close to one without losing all the configurations. The repository contains an edited transcription of this paper and a collection of local quantitative refinements [192]. We use the transcription's explicit largeness hypothesis for Lemma 15.5, rather than silently suppressing it.

The local question is concrete:

Is the extremely large exponent produced by repeating the degree-one cosine kernel intrinsic to this arrangement-selection problem?

We answer this question negatively under an explicit, substantially smaller field-size threshold. This is a question extracted from a proof and its constants; we do not assert that it was previously advertised as a named open problem.

There are two main changes. First, the full family of squarefree coordinate monomials can be replaced by the single monomial consisting of the product of all $k + 1$ coordinates. The relevant requirement is not linear independence over the whole finite field, which would be impossible for so many scalar values. It is the absence of *short* unintended relations, and a polynomial zero-count controls these. Second, a high-degree Fejér kernel has an arrangement survival advantage proportional to its bandwidth, rather than an advantage $1 + 2^{1-s}$ barely exceeding one. This makes one selection sufficient.

F.45.2 Main explicit statement

Definitions of arrangements and respect are given in Section II.165. All arrangements are ordered and counted with their parameters; zero sidelengths and repeated vertices are allowed in the counts.

Theorem F.45.1 (Explicit one-stage selection). *Let $k, d \geq 1$, let $s = 2d2^k$ and $M = (2d+1)k+2d-1$, and let $0 < \alpha, \beta, \eta \leq 1$. Define*

$$L = \left\lceil \frac{2(s+1)}{\alpha\eta} \right\rceil, \quad \Gamma_s(\alpha, \eta) = \frac{\alpha^s \eta^{s-1}}{3^{s-1}(s+1)^s}. \quad (\text{F.45.1})$$

Let q be a prime power of characteristic p satisfying

$$p > 2(L-1), \quad q \geq (k+1)\beta^{-(2d-1)} \left(\frac{18(s+1)^2}{(\alpha\eta)^2} \right)^s. \quad (\text{F.45.2})$$

Suppose $B \subseteq \mathbb{F}_q^{k+1}$ has $|B| = \beta q^{k+1}$ and $\phi : B \rightarrow \mathbb{F}_q$ respects at least $\alpha\beta^{2d-1}q^M$ d -arrangements in B . Then there is $B' \subseteq B$ containing at least

$$\Gamma_s(\alpha, \eta) \beta^{2d-1} q^M \quad (\text{F.45.3})$$

d -arrangements. In fact at least this many are respected, and the proportion of all arrangements in B' that are respected is at least $1/(1+\eta) \geq 1-\eta$. Moreover,

$$\frac{|B'|}{q^{k+1}} \geq \beta \Gamma_s(\alpha, \eta)^{1/(2d-1)}. \quad (\text{F.45.4})$$

If $q = p$ is prime, the field-size inequality in (F.45.2) already implies its characteristic inequality.

The theorem is proved in Section II.169. The more flexible Theorem II.169.2 uses the exact finite Fejér sum and a smaller exceptional-set constant; (F.45.2) is a convenient closed-form sufficient condition, not an optimized one.

F.45.3 Scope and relationship to existing work

The basic probabilistic architecture comes from Gowers’ Lemma 15.5. The features, bandwidth, exceptional-set accounting, and one-stage parameter choice are reworked here. The algebraic and analytic tools are elementary, and the note supplies their proofs to make the quantitative conclusions independently auditable.

A local improvement should not be confused with a new global estimate for progression-free sets. For perspective, Leng, Sah, and Sawhney [2] prove, for each progression length $\ell \geq 5$, a bound $r_\ell(N) \ll N \exp(-(\log \log N)^{c_\ell})$ using a different package of inverse-theorem and density-increment inputs. Nothing here improves or competes with that theorem. Propagating the present constants through all later steps of Gowers’ argument is a separate task.

The repository comparison is scoped to the Section 15 source and the local-refinement index inspected for this note. A literature search is not a proof of priority. The results below are presented as proved refinements for review and integration, rather than as a certified first discovery or a claimed solution of a major global open problem.

[write] **7 October 2026; sources 40–46.** *Identifications within this report.* [192] is this repository: “the repository’s edited statement” of Lemma 15.5 is the transcription’s Section 15, and the “local-refinement index” is this directory’s README; source 40 records no commit or blob, only an inspection date. Source 40 was placed by 2dfe46f44, after source 35, which builds on it and cites it as [160] (note in Section F.40): that forward reference is now resolved, and source 35’s degree-weighted selection (Section II.144) improves source 40’s field-size exponent (note after Theorem II.169.2). Earlier Fejér-kernel material in this report: source 17’s Fejér filters at Section 9 (Section II.57); the selection theorems of sources 25 and 33 are neighbours. The intake re-derived the proofs and reproduced the scale table of Section II.170 to the printed digits; source 40’s suite was rerun on a copy at the write, with identical output.

F.46 Source 41. The theorem and its scope

Source 41, delivered Section 1.

Added in the eighth write, 7 October 2026. Delivered title block. “Threshold free polynomial density increments on cyclic groups / Report277 / Research report prepared for Vladimir / 6 October 2026”; *the PDF author field reads “Research report prepared for Vladimir”; no statement on AI assistance is made.* Archive Report277_Threshold_Free_All_Degree_Density_Increments.zip (arrival a6eea2a19), manuscript article.tex (1,022 lines, 21 pages), pin c94998a5d (its SOURCES.md, shipped as 41-all-degree-increments-SOURCES.md). It is the third report of the series of sources 26 (Report275) and 32 (Report276), and its Sections 2 and 5 are source 32’s Sections 2 and 4 word for word (note in Section I.63). Printed below: the abstract and Section 1; Sections 2–7 form, with source 42, a chapter of Part I (Section I.62), Sections 8–9 are in Part V. Its macro \norm omits \left/\right, as source 26’s did, and prints with this report’s.

Abstract of source 41 (as delivered). For every polynomial degree $k \geq 1$ and every cyclic modulus, a disjoint comparable cover carrying total balanced polynomial correlation αN yields a proper progression of density at least $\delta + \alpha/3$ and length at least

$$\max\{1, (N/M)^{1/D_k}/8\}, \quad D_1 = 4, \quad D_k = \frac{1}{4}(k!)^2 2^{(k+1)^2} + 1 \quad (k \geq 2).$$

There is no missing large-scale premise. We prove every analytic input, including a coefficient-uniform monomial recurrence estimate with explicit constants, the sharper quadratic base, a lossless recursive phase partition, and a subgroup-coset variance alternative over arbitrary cyclic groups. We also retain the exact recursive partition budget and sharper affine, quadratic, and cubic alternatives. For $k \geq 2$, a compact bound gives length at least $\max\{1, 2^{-(k+8)}(N/M)^{1/(12\mathcal{P}_k+1)}\}$, where $\mathcal{P}_k = k!2^{k(k+1)/2}$. The conclusion implies the entire pinned Corollary 5.8 interface in every degree. This is a written mathematical theorem, not a Lean verification, a global quantitative Ramsey or Szemerédi improvement, or a priority claim. The companion provides bounded exact diagnostics and reproducible packaging; finite samples do not prove recurrence for arbitrary real coefficients.

For a real t , write $e(t) = \exp(2\pi it)$ and $\|t\| = \min_{v \in \mathbb{Z}} |t - v|$. Let $G = \mathbb{Z}/N\mathbb{Z}$, where $N \geq 2$ is an integer, and put $e_N(y) = e(\tilde{y}/N)$ using any integer representative of $y \in G$. A *proper modular arithmetic progression* is a set

$$P = \{u + jd : 0 \leq j < n\} \subseteq G$$

whose displayed parameterization is injective. Its length is n . Modular wrapping is allowed; the step need not be a unit. A singleton is proper. The polynomials considered here are *coefficient polynomials*

$$\phi(x) = \sum_{j=0}^k a_j x^j, \quad a_j \in G.$$

This includes lower-degree and constant polynomials. Over composite moduli it does not silently substitute an arbitrary finite-difference definition of polynomial.

Let $A \subseteq G$, $\delta = |A|/N$, and $f = \mathbf{1}_A - \delta$. Define

$$K_k = (k!)^2 2^{(k+1)^2}, \quad D_1 = 4, \quad D_k = K_k/4 + 1 \quad (k \geq 2). \quad (\text{F.46.1})$$

These are positive integers. In particular $K_1 = 16$ and $K_2 = 2048$.

Theorem F.46.1 (All degree local cover increment). *Let $k, M \geq 1$ be integers. Let $P_1, \dots, P_M$ be nonempty pairwise disjoint proper modular progressions such that*

$$A \subseteq U := \bigcup_{i=1}^M P_i, \quad |P_i| \leq 2|P_j| \quad (1 \leq i, j \leq M). \quad (\text{F.46.2})$$

Let each ϕ_i be a coefficient polynomial of degree at most k . Suppose $\alpha > 0$ and

$$\sum_{i=1}^M \left| \sum_{x \in P_i} f(x) e_N(-\phi_i(x)) \right| \geq \alpha N. \quad (\text{F.46.3})$$

There is a nonempty proper modular progression Q such that

$$|Q| \geq \max \left\{ 1, \frac{(N/M)^{1/D_k}}{8} \right\}, \quad |A \cap Q| \geq (\delta + \alpha/3)|Q|. \quad (\text{F.46.4})$$

There is no lower bound on α or N/M , and no primality or least-prime-factor hypothesis. The conclusion does not require $Q \subseteq U$.

The construction has two regimes. If the correlation is small at the chosen integer target length, a second moment gives the requested density increment directly. Otherwise the correlation forces enough mass in the cover that every parent meets the explicit phase-partition threshold. Thus “threshold free” describes the final hypotheses, not the absence of internal quantitative estimates.

The proof is self-contained. Section I.63 establishes the quadratic base and explains the Fourier detection step in full. Section I.64 proves recurrence in every degree with an explicit ordered-divisor constant. Section I.65 constructs the partitions and verifies their all-degree budget. Sections I.66 and I.67 prove the moment alternative and final assembly. Section I.68 preserves the actual recursive budget and sharper small-degree conclusions. Section V.154 compares the exact pinned source interface and credits the classical mechanisms.

Throughout, k is the polynomial degree. In repeated differencing the number of shifts is $s = k - 1$, and the power on an exponential sum is $w = 2^s$. These are separate from the density exponent D_k . The recurrence accuracy R is an integer at least two; the progression target length is L , also an integer. Locally introduced H, T are explicitly integer scales. All logarithms without a subscript are natural.

[write] **7 October 2026; sources 40–46.** *Identifications and credit within this report.* [193] is this repository at the pin c94998a5d (the predicate `corollary_5_8` and the definition of `Polynomial0n`). Source 41 is Report277 of the series whose Report275 and Report276 are sources 26 and 32; *its article cites neither*, although it is an extended edition of source 32. Its Sections 2 and 5 are source 32’s Sections 2 and 4 word for word (Sections I.45 and I.47; notes there), its Corollaries I.68.4 and I.68.5 retain source 26’s Corollary I.41.3 and source 32’s Theorem F.37.1, and the two-regime method (variance for small α , long parents for large α) is theirs; the intake’s 8-word shingles find 47.9 % of source 32’s text and 4.4 % of source 26’s in source 41. Source 41’s own `SOURCES.md` names sources 26 and 32 only to correct the SHA-256 values in their staged `SOURCES.md` files, which hash each file followed by one line feed (checked: the four raw values match). What is new here is the recurrence in every degree (Theorem I.64.3), the recursive localization with its budget (Theorem I.65.3, Lemma I.65.4), the theorem in every degree, the compact and limiting budgets, and the cubic bound $(N/M)^{1/4096}/8$.

It implies the formal statement (checked at the write). `corollary_5_8` (`Section05.lean`, line 242, unchanged since source 41’s pin) asks for a proper Q with $(N/M)^{1/K_k}/8 \leq |Q|$ and $(\delta + \alpha/16)|Q| \leq |A \cap Q|$, $K_k = \text{polynomialPartitionConstant } k = (k!)^2 2^{(k+1)^2}$. Since $D_1 = 4 \leq 16 = K_1$ and $D_k = K_k/4 + 1 \leq K_k$ for $k \geq 2$, and $N/M \geq 1$, the length $(N/M)^{1/D_k}/8$ dominates; $\alpha/3 \geq \alpha/16$; the empty-parent and $N = 1$ cases are vacuous, as source 41 says in Section V.154. So Theorem F.46.1 implies the formal conclusion in every degree, on paper. *The formal statement is now exact* at HEAD 6d4942222 (ledger 113 exact companions and 7 open statements): `corollary_5_8_holds` (`Proofs05ThresholdFreeLocalization.lean`, line 86) was committed in b548d998e while this write was in progress, by a formal proof with the two-regime architecture of sources 26, 32 and 41 and the repository’s own finite localization (Section I.62). Source 41’s own theorem, with $1/D_k$ and $\alpha/3$, and its recurrence are not formalized.

F.47 Source 42. The arithmetic increment and its meaning

Source 42, delivered Section 1.

Added in the eighth write, 7 October 2026. Delivered title block. “Sharp quadratic scale density increments on cyclic groups / Report278 / Research report prepared for Vladimir / 7 October 2026”; *the PDF author field reads “Research report prepared for Vladimir”; no statement on AI assistance is made.* The title page is dated 7 October 2026; the archive arrived at 17:40 on 6 October 2026 (Pacific time). Archive `Report278_Sharp_Quadratic_Scale_Density_Increments.zip` (arrival a6eea2a19), manuscript `article.tex` (688 lines, 15 pages), pin b16ce380d for this report’s article, with sources 41 and 26 identified by SHA-256 (its `SOURCES.md`, shipped as `42-quadratic-scale-SOURCES.md`). Printed below: the abstract and Section 1; Sections 2–7 follow source 41 in the chapter of Part I (Section I.62), Sections 8–9 are in Part V.

Abstract of source 42 (as delivered). Let $N \geq L \geq 2$, let $A \subset \mathbb{Z}/N\mathbb{Z}$ have density $0 < \delta < 1$, and set $m = \max\{h : h \mid N, h < L\}$ and $q = N/m$. There is a proper L -term modular arithmetic progression of density at least

$$\delta + \frac{1 - \delta}{L} \frac{\max\{q - L, 0\}}{q - 1}.$$

An independent-difference step distribution makes every off-diagonal covariance nonnegative while excluding every nonzero short-order step. Its support size and zero-step mass are exactly optimal in natural classes; the stronger positive-definite weight statement is an instance of the classical Turán packing bound. In particular, the gain $(1 - \delta)/(2L)$ holds when $N \geq 2(L - 1)^2 + 1$. For each real $C > 1$, the best constant uniform over all nontrivial densities under $N \geq CL^2$ is $1 - 1/C$ in front of $(1 - \delta)/L$. Prime-quotient sets attain the arithmetic formula exactly. A separate sliding-window argument gives gain greater than $\delta(1 - \delta)/(2L)$ when $N > L(L - 1)$, with a sharp strict threshold even at densities bounded away from zero and one. We also prove finite-abelian and bounded-function versions and a sparse affine-cover consequence with a square-root modulus cap. The proofs are self-contained; the accompanying exact computations are bounded diagnostics, not formal certificates or substitutes for proof. No literature-priority or global quantitative Szemerédi or Ramsey claim is made.

Let $G = \mathbb{Z}/N\mathbb{Z}$ with its additive group law. A *proper modular arithmetic progression* of length L is

$$P(a, d; L) = \{a + jd : 0 \leq j < L\}$$

with an injective displayed parameterization. Equivalently, the additive order of d is at least L . Modular wrapping is allowed, and d need not be a unit modulo N . A singleton is also called proper. All group averages in this report are normalized.

For $N \geq L \geq 2$ define

$$m = m(N, L) := \max\{h \in \mathbb{Z}_{\geq 1} : h \mid N, h < L\}, \quad q = q(N, L) := \frac{N}{m}. \quad (\text{F.47.1})$$

The set defining m contains 1. Since $m < L \leq N$, one has $q \geq 2$. For a nonempty proper subset A of G , put

$$\delta = \frac{|A|}{N}, \quad M_L(A) = \max_{\text{ord}(d) \geq L} \frac{1}{L} \sum_{j=0}^{L-1} \mathbf{1}_A(a + jd), \quad H_L(A) = M_L(A) - \delta,$$

where the maximum also ranges over $a \in G$. Proper steps exist because the step 1 has order N .

Theorem F.47.1 (Arithmetic increment). *If $N \geq L \geq 2$ and $0 < \delta < 1$, then*

$$H_L(A) \geq \frac{1 - \delta}{L} \frac{\max\{q - L, 0\}}{q - 1}. \quad (\text{F.47.2})$$

When $q > L$, a progression giving the displayed gain can be chosen with step represented by an integer d such that $0 < |d| < q$.

The exact arithmetic parameter matters. If N is prime, then $m = 1$ and $q = N$, so the coefficient is $(N - L)/(N - 1)$. In contrast, when $N = q(L - 1)$ with $q \geq L$ prime, the largest divisor below L is $L - 1$. In this second family the theorem is attained by sets of density $1/q$. Thus a large modulus alone is not a substitute for tracking its short-order directions.

The useful size-only conclusions have a simpler form.

Corollary F.47.2 (Quadratic scale). *Let $D = \lfloor (N - 1)/(L - 1) \rfloor$. Then*

$$H_L(A) \geq \frac{1 - \delta}{L} \max \left\{ 1 - \frac{L - 1}{D}, 0 \right\}. \quad (\text{F.47.3})$$

In particular:

- (i) $N > L(L - 1)$ guarantees a strictly positive increment;
- (ii) $N \geq 2(L - 1)^2 + 1$ guarantees $H_L(A) \geq (1 - \delta)/(2L)$;
- (iii) for every real $C > 1$, $N \geq CL^2$ guarantees $H_L(A) \geq (1 - 1/C)(1 - \delta)/L$.

The coefficient in (iii) is best possible uniformly over N, L, A satisfying that size condition. This optimality includes densities tending to zero; it is not an optimality assertion at a fixed positive density.

[write] **7 October 2026; sources 40–46.** Since $2(L - 1)^2 + 1 < L^3$ for every $L \geq 2$, part (ii) weakens the hypothesis of the variance lemma of sources 26, 32 and 41 (Lemma I.66.1); that lemma, with $L^3 \leq N$, is the formal `exists_proper_progression_density_increment` (`Proofs05VarianceIncrement.lean`, line 89, at HEAD 6d4942222). Source 41’s exponents do not change: its high branch needs $N/M > (4L)^{D_k}$, far above L^3 .

Although the density-uniform coefficient degenerates at the quadratic endpoint, Section I.73 proves a density-sensitive gain there. In particular, $N > L(L - 1)$ gives $H_L(A) > \delta(1 - \delta)/(2L)$. This is compatible with the sparse equality family and provides a positive coefficient when density is confined away from zero and one.

The standalone variance lemma in Report277 [196] gave the gain in (ii) under $N \geq L^3$. The present quadratic sufficient condition genuinely strengthens that lemma. It does not change the separate phase-localization exponents in Report277. Section I.74 gives a precise consequence where the improved variance scale changes a modulus-dependent cap.

The argument has three ingredients. Sampling a step as the difference of independent identically distributed variables gives a square identity for every covariance, including multipliers that are not units. An arithmetic interval makes every nonzero sampled step proper. Finally, removing the zero step and applying the elementary variance–range inequality gives a one-sided density increment. Sections I.69–I.71 separate these ingredients and their exact optimization; Section I.72 establishes the sharpness claims.

[write] **7 October 2026; sources 40–46.** *Identifications within this report.* [196] (Report277) is source 41 and [195] (Report275) is source 26, identified by the SHA-256 of their delivered manuscripts (both match); [194] is this report at `b16ce380d`, its third write (sources 01–14), whose lines 4804–4846 hold source 10’s variance-to-density calculation, credited. “The standalone variance lemma in Report277” is source 41’s Lemma I.66.1, which is also source 26’s Lemma I.42.1 and source 32’s Lemma I.47.1. Source 42 does not cite source 32. The intake checked the arithmetic bound by brute force on 2,526 pairs (A, L) with $N \leq 48$ (72 equalities, no violation); at the write the prime-quotient equality $H_L(A) = (p - L)/(Lp)$ was recomputed in exact rationals for $p = 5, 7, 11, 13$ and every $2 \leq L \leq p$.

F.48 Source 43. The research problem and the principal results

Source 43, delivered Section 1.

Added in the eighth write, 7 October 2026. Delivered title block. “Sharp Boolean Phase Integration from Gowers Derivative Spectra / Exact cubic obstruction ranks, repeated-variable certification, finite-abelian orthogonality, and a quartic reduction / Research manuscript prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt project / 6 October 2026”; *the PDF author field reads “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov”*. Archive `Boolean_Phase_Integration_Research.zip` (arrival 7382d8c44), manuscript `boolean_phase_integration.tex` (1,938 lines, 26 pages), pin 797aa3cfe (its `data/source_audit.json`, shipped as `data/43-boolean-phase-source_audit.json`). Printed below: the abstract, the integration note of its title page and Section 1; Sections 2–9 follow source 46 in a chapter of Part IV (Section IV.25), Sections 10–13 are in Part V. Its labels already carry the stem `bpi:`, kept under this report’s prefix. Its macro `\shortpin` is renamed `\shortpinfortythree`; its ‘question’ environment (“Research question”, numbered 1–10 on its own counter) is set in this report’s ‘Research question’ style, numbered within its section.

Abstract of source 43 (as delivered). A multilinear frequency pattern with large derivative Fourier energy need not integrate to a polynomial phase in small characteristic, even when it is symmetric. We study the exact obstruction for one bounded function. For a symmetric trilinear form T over $\mathbb{F}_2$, put $\Omega_T(a, b) = T(a, a, b) + T(a, b, b)$. We prove that the largest possible tested energy is exactly $2^{-\text{rank } \Omega_T/2}$, and that $\text{rank } \Omega_T/2$ is also the least codimension of a subspace on which T has a nonclassical cubic primitive. Explicit extremizers establish both assertions sharply. In every degree $d \geq 2$, energy greater than $1 - 2^{1-d}$ certifies a nonclassical primitive and gives lossless phase removal; this threshold is sharp. The additional repeated-variable identities are detected by an exact shear identity and an alternating-rank profile. A division-free orthogonality theorem treats every finite abelian group, including characteristic two and composite moduli. Its mixed version improves the exponent eight in a related bilinear correlation estimate to the optimal exponent four, and gives a codimension $2\log_2(1/\delta)$ refinement of Tidor’s symmetric trilinear repair lemma. Finally, the nonintegrable symmetric quartic extremum is $11/16$ in dimension two; in all dimensions the remaining extremal problem reduces to one explicit two-coordinate tensor. The proofs are given in full. The underlying nonclassical integration theorem and finite Heisenberg algebra are established background, not claimed discoveries. These are local research results, not a new end-to-end Szemerédi bound; priority and proof-assistant verification have not been established.

Integration note. The comparison with the formal project uses snapshot 797aa3cfe967. Labels in this manuscript have prefix `bpi:`. The package contains exact finite checks, a source audit, and a proposed formalization interface; it does not change the repository’s checked theorem count.

F.48.1 From a selected Fourier coefficient to an integrable tensor

Gowers’s phase-removal step in Section 17 of [1] converts a coherent family of derivative Fourier coefficients into polynomial structure. Proposition 17.2 is a global model for the localized argument that follows. The corrected formal statement in ProveIt explicitly uses degree $k + 1$ and factorial invertibility [198]. The question here is what survives when the ambient group is a vector space in characteristic two, where division by that factorial is unavailable and symmetry is not the entire integrability condition.

Throughout the Boolean parts, $V = \mathbb{F}_2^n$, all averages are uniform and normalized, and

$$\partial_h f(x) = f(x + h) \overline{f(x)}, \quad \hat{g}(\xi) = \mathbb{E}_x g(x) (-1)^{\xi(x)}. \quad (\text{F.48.1})$$

For a d -linear form $T : V^d \rightarrow \mathbb{F}_2$, $d \geq 2$, define

$$\mathcal{E}_T(f) = \mathbb{E}_{h_1, \dots, h_{d-1}} \left| \mathbb{E}_x \partial_{h_1} \cdots \partial_{h_{d-1}} f(x) (-1)^{T(h_1, \dots, h_{d-1}, x)} \right|^2. \quad (\text{F.48.2})$$

Thus T includes the frequency-evaluation variable, not just the derivative directions. We assume $\|f\|_\infty \leq 1$ unless an amplitude-sensitive statement says otherwise.

A *nonclassical polynomial* of degree at most d is a function $P : V \rightarrow \mathbb{T} = \mathbb{R}/\mathbb{Z}$ whose $(d+1)$ -fold additive differences vanish. A primitive of T satisfies

$$\Delta_{h_1} \cdots \Delta_{h_d} P(x) = \tfrac{1}{2} T(h_1, \dots, h_d) \quad \text{in } \mathbb{T}. \quad (\text{F.48.3})$$

Here $t \mapsto t/2 \pmod{1}$ is the additive embedding of $\mathbb{F}_2$ in $\mathbb{T}$. The value group is essential; a primitive taking values only in $\{0, 1/2\}$ generally does not exist.

F.48.2 An exact cubic invariant

For a symmetric trilinear T set

$$\Omega_T(a, b) = T(a, a, b) + T(a, b, b), \quad \mathcal{M}(T) = \max_{\|f\|_\infty \leq 1} \mathcal{E}_T(f). \quad (\text{F.48.4})$$

The maximum exists because the domain is finite and the closed unit polydisc is compact. Our central conclusion is

$$-\log_2 \mathcal{M}(T) = \frac{1}{2} \text{rank } \Omega_T = \min\{\text{codim } U : T|_{U^3} \text{ integrates}\}. \quad (\text{F.48.5})$$

The rank is the rank of the alternating bilinear form Ω_T , not the partition rank of T . Section IV.39 proves the formula, gives normal forms and extremizers, and retains the stronger amplitude factor $\| |f|^2 \|_{U^2}^4$.

In particular, $\mathcal{E}_T(f) > 1/2$ forces a symmetric cubic tensor to integrate. At energy $1/2$ there is a genuine nonintegrable example. At smaller energies, the exact logarithmic codimension law remains informative:

$$\mathcal{E}_T(f) \geq a > 0 \implies \text{an integrable restriction of codimension at most } \lfloor \log_2(1/a) \rfloor. \quad (\text{F.48.6})$$

This is a statement about a prescribed tensor. It is more specific than merely correlating f with some polynomial phase.

F.48.3 All-degree certification and a first quartic obstruction

The cubic result sits inside a general certification theorem:

$$\mathcal{E}_T(f) > 1 - 2^{1-d} \implies \exists P \text{ satisfying (F.48.3),} \quad \mathcal{E}_T(f) = \|f e(-P)\|_{U^d}^{2^d}, \quad (\text{F.48.7})$$

where $e(t) = \exp(2\pi i t)$. The threshold is sharp in every dimension at least two. No symmetry hypothesis on T is needed in (F.48.7); it is first forced by the energy.

Once symmetry is already known, a stronger estimate detects the remaining obstruction. For $d \geq 3$ define

$$\Omega_z(a, b) = T(a, a, b, z) + T(a, b, b, z), \quad z \in V^{d-3}. \quad (\text{F.48.8})$$

Then

$$\mathcal{E}_T(f) \leq \mathbb{E}_z 2^{-\text{rank } \Omega_z / 2}. \quad (\text{F.48.9})$$

In particular a nonzero repeated-variable obstruction forces $\mathcal{E}_T(f) \leq 1 - 2^{2-d}$. This bound is sharp for $d = 3$. Its optimal higher-degree replacement is a separate question.

For $d = 4$, we give more than an interval for that question. Every nonintegrable symmetric quartic with energy greater than $5/8$ is, up to an integrable summand and linear coordinates, the pullback of

$$C_4(h_1, h_2, h_3, h_4) = \sum_{i=1}^4 v(h_i) \prod_{j \neq i} u(h_j) \quad (\text{F.48.10})$$

for two independent linear forms u, v . In dimension two its exact maximum energy is $11/16$. In larger dimensions the remaining issue is whether extra directions on which the tensor is zero can increase this maximum. Section IV.41 proves the reduction and the exact two-dimensional theorem.

F.48.4 What is strengthened, and what is not claimed

The preceding research manuscript [197] established an odd-characteristic symmetry threshold and explicitly asked for characteristic-two, nonclassical-certification, and composite-modulus extensions. The present results answer its characteristic-two question, its nonclassical-certification question over $\mathbb{F}_2$ in all degrees, and its finite-abelian bilinear/rank-profile question. They do not answer nonclassical certification in every odd small characteristic.

The algebraic characterization and integration of nonclassical symmetric multilinear forms are existing results; see Tao–Ziegler [184] and Tidor [150]. We give a direct Boolean construction so the new energy arguments have a complete, explicit endpoint. Likewise, the finite Heisenberg representation mechanism has established antecedents [185]; it is not presented as a new representation theory.

The proposed contributions are the exact derivative-spectrum obstruction laws, their sharp phase-removal consequences, the normalized finite-abelian estimates, and the refinements developed from them. In Section IV.40, the bilinear exponent and the symmetric trilinear codimension comparison are made under the actual published hypotheses, rather than by substituting a one-function theorem into a seven-function problem. A targeted source comparison is not a proof of priority. None of these results supplies the missing global selector, handles the boundary of an arithmetic box, or changes a final Szemerédi bound without further work.

[write] 7 October 2026; sources 40–46. *Identifications and credit within this report.* [197] is source 36 (Section F.41), read from the user’s library; its questions `ssc:q:two`, `ssc:q:nonclassical` and `ssc:q:composite` are Research questions V.136.9, V.136.7 and V.136.8 here (item IV.9). [198] is this repository at `797aa3cfe`. *Source 46, of the same day, proves the main theorem independently* (its Theorem IV.27.1, eight functions) and the same codimension $2 \log_2(1/\delta)$ (its Theorem IV.28.1); neither cites the other, and their texts share 0.6% (intake). Source 46 is printed first in the chapter of Part IV (Section IV.25); this source is a second route there, with its further results.

The four claims the intake did not check (checked at the write; script `check43.py` of the write record). (a) Theorem IV.35.2: the bound held for random g, h and eleven biadditive forms on groups including $(\mathbb{Z}/2\mathbb{Z})^3$, $\mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/4\mathbb{Z}$, $(\mathbb{Z}/4\mathbb{Z})^2$, $(\mathbb{Z}/6\mathbb{Z})^2$, $\mathbb{Z}/8\mathbb{Z} \times \mathbb{Z}/4\mathbb{Z}$ and $(\mathbb{Z}/12\mathbb{Z})^2$ (largest ratio 0.988); d_B was an integer in every case; the proof’s construction (a self-annihilating H of index d_B and a joint eigenvector of unit modulus on each coset) attained $1/d_B$ exactly in all nine cases with $|G| \leq 36$; the Smith-form example on $(\mathbb{Z}/12\mathbb{Z})^2$ was reproduced for $s = 1, 2, 3, 4, 6$. (b) Theorem IV.38.4: on $\mathbb{F}_2^2$, over all 16 bilinear and all 256 trilinear forms and all 4,096 functions with values among the eighth roots of unity, the energy of forms outside nCSM^d never exceeded $1 - 2^{1-d}$ and attained it, symmetric ones outside nCSM^3 stayed at most $1/2$, and a search over all candidate primitives confirmed Theorem IV.37.2 (8 integrable forms for $d = 2$ and $d = 3$); the extremizer of degree four has energy exactly $7/8$. (c) Theorem IV.40.1: the bias identity held exactly on six forms,

alternating maximization over unit vectors reached $d_B^{-1/2}$ and never exceeded it, and the $\mathbb{F}_2^2$ extremizer gives $|L|^4 = 1/4$ exactly; Tidor’s Lemma 4.1, re-read in the arXiv HTML of version 2, concludes the bias bound δ^8 , so the fourth power is a strict local refinement of that display, as stated. (d) Theorem IV.41.3: the 29-monomial identity agreed with the energy to 10^{-15} for random complex f ; $\mathcal{E}_{C_4}(1) = 11/16$; the cosine bound 56 is attained; random ascent found 0.6874999; all 24 symmetric quartics on $\mathbb{F}_2^2$ outside nCSM⁴ have maximum exactly 11/16 over the sixteenth roots of unity. The proofs were read; the finite checks are evidence for them, not replacements.

F.49 Source 44. The research target and what is new

Source 44, delivered Section 1.

Added in the eighth write, 7 October 2026. Delivered title page. “Additive combinatorics / finite geometry / Finite-Grid Formulae and Exact Core-Avoidance Laws for Gowers Arrangements / Computing the second degeneracy coefficient in polynomially many arithmetic operations / An AI-assisted research manuscript prepared for Vladimir Reshetnikov and the ProveIt repository / 6 October 2026”; *the PDF author field reads “AI-assisted research manuscript prepared for Vladimir Reshetnikov”*. Archive `gowers_grid_core_research.zip` (arrival `0b4890a1c`), manuscript `article.tex` (1,632 lines, 24 pages), pin `8c12c906e` (its `SOURCE_AUDIT.md`, shipped as `44-grid-core-SOURCE_AUDIT.md`). Printed below: the abstract, the status statement of its title page and Section 1; Sections 2–10 form a chapter of Part II (Section II.173), Sections 11–13 and Appendix A are in Part V. Its macros `\T` (the set $\{-1, 0, 1\}$), `\card` (an absolute value) and `\sha` (its pin) are renamed `\Tfortyfour`, `\cardtwentyfive` (source 25’s macro of the same definition) and `\shafortyfour`; its ‘question’ environment (“Question”, on the shared counter) is set in this report’s ‘Question’ style. Its “Baseline A” and “Baseline B” are sources 25 and 27 (note in its Section 1).

Abstract of source 44 (as delivered). We resolve two explicit questions in the recent ProveIt research on Gowers’s arrangement-degeneracy lemma: computation of the cross-section intersection statistic in arbitrary dimension, and classification of the intermediate-characteristic core directions. The second Whitney coefficient $b_{2,d}$ of the balanced ternary cross-section arrangement is expressed through subsets of the fixed nine-point grid $\{-1, 0, 1\}^2$. The resulting algorithm takes $O(d^3)$ exact arithmetic operations, with a constant independent of d . A finite rational certificate compresses the answer to eight elementary lattice-walk energies. In particular,

$$b_{2,8} = 3\,280\,234\,631\,541.$$

This supplies the previously unevaluated second coefficient in the 8-arrangement case used in Section 15. A determinant argument makes that coefficient characteristic-independent for $p > \max(2d, 16)$.

We also count the union of all core events exactly, replacing pairwise inclusion–exclusion by a product. For every characteristic $p \geq 5$ the conditional avoidance probability is $\prod_{i < k} (1 - (2^{2d} - 1 + 2i)/q^2)$; in characteristics two and three it is $\prod_{i < k} (1 - p^{2d-1+i}/q^2)$. The width classification requires no hypothesis involving k . These laws give contrasting growing-dimension limits and universal degeneracy thresholds. Finally, the deficit of $b_{2,d}$ from $\binom{H_d}{2}$ has an explicit triangular-lattice asymptotic of order $49^d/d$.

Status. Conventional proofs, with a separately identified finite rational certificate for the eight-shape compression. Executed exact checks are supplied. No independent referee review, proof-assistant verification, or literature-priority certification is claimed. These are local strengthenings of the degeneracy analysis, not a new global quantitative Szemerédi theorem.

Gowers’s Lemma 15.4 controls exceptional labelled arrangements before the Riesz-product selection in Lemma 15.5. Its prime-field bound, in the notation below, is $k3^{2d^k}/q$ for the exceptional

proportion [1, Section 15, pp. 562–565]. The full definition tests every squarefree coordinate moment, including those containing the final coordinate. We retain that definition throughout.

Two earlier manuscripts are important baselines. The first classifies the codimension-one obstructions [199]. The second classifies the core, or second-order, coefficient directions, gives a cubic remainder, and expresses the second asymptotic coefficient using a cross-section intersection invariant $b_{2,d,p}$ [200]. In particular, the core count $k(2^{2d} - 1) + k(k - 1)$ and the second coefficient for $d = 2$ are *not* claimed as new here. The latter manuscript explicitly asks for an efficient exact algorithm for $b_{2,d,p}$, especially for $d = 8$, and for the core classification when $5 \leq p \leq 2k + 2$.

The main advances of this manuscript are as follows. The fixed-grid formula in Theorem II.179.1 solves the first problem in the stable characteristic regime for every d , rather than merely computing a few small examples. Lemma II.175.1 solves the second problem for every $p \geq 5$, including all intermediate characteristics. Theorems II.176.2 and II.176.3 count the complete core union, not just its first two inclusion–exclusion terms. Theorem II.178.1 supplies a dimension-independent determinant cutoff for the relevant finite-field incidence structure. Theorem II.181.2 gives a second, exponentially smaller scale in the growth of the cross-section statistic.

F.49.1 A headline specialization

Let $Z_{k,d}(q)$ count degenerate labelled d -arrangements of k -cubes, and put $D = (2d + 1)k + 2d - 1$. For every fixed $k \geq 1$, uniformly over finite fields of characteristic greater than 16, the results below give

$$Z_{k,8}(q) = (k + 2\,585\,444)q^{17k+14} \tag{F.49.1}$$

$$+ \left\{ \binom{k}{2} - 2\,519\,909k - 3\,280\,234\,631\,541 \right\} q^{17k+13} + O_k(q^{17k+12}). \tag{F.49.2}$$

The second coefficient is exact, including its negative sign. It is not obtained by fitting counts at a few prime powers. The large negative contribution records overlaps among the cross-section hyperplanes.

The statement is a fixed- k asymptotic with a uniform characteristic range. It is not an estimate with a small implied constant uniformly in k . The explicit remainder in Theorem II.177.2 still contains a large inventory of non-core directions.

F.49.2 Source and novelty boundaries

The inspected repository interface is `Combinatorics/Ramsey/Lean/GowersSzemerédi/Sections14_15.lean` at commit `8c12c906e4ab5a00001a65de3c882f935a15661a` [201]. We follow its modular-multiple exclusion, not an exclusion of only two integer arrays. The exact results in the two baseline manuscripts are attributed there, and the argument here reproves the finite-dimensional ingredients needed to use them.

The finite-field method and Whitney coefficients of hyperplane arrangements are classical [98, 99]. Related mid-hyperplane arrangements also encode equal subset sums [188]; their existence does not identify their characteristic polynomials with the balanced restriction considered here. Our contribution is the specific fixed-grid reduction and its consequences. A targeted source comparison does not establish historical priority for every individual identity.

[write] **7 October 2026; sources 40–46.** *Identifications within this report.* [199] (“Baseline A”) is source 25 (Section F.30) and [200] (“Baseline B”) is source 27 (Section F.32), cited here through the file names

of the user’s library, not by repository path (the titles match; the identifiers in `SOURCE_AUDIT.md`, shipped as `44-grid-core-SOURCE_AUDIT.md`, are recorded, not edited). [201] is this repository at `8c12c906e`. The questions answered are source 27’s on the algorithm for $b_{2,d,p}$ and on intermediate characteristics (item II.38); the core count and the $d = 2$ coefficient are source 27’s (Theorem II.97.6, Theorem F.32.3), the leading term source 18’s, as source 44 says. The intake reproduced $b_{2,d} = 3, 889, 97,081$ for $d = 2, 3, 4$ by brute force over pairs of hyperplanes, and $b_{2,d,p}$ equals the rational value already for $p = 7, \dots, 23$ at $d = 3$ and $p = 11, \dots, 17$ at $d = 4$, below the cutoff 16 (a data point for its question on sharp exceptional characteristics, Section V.163); source 44’s suite was rerun on a copy at the write, with identical output.

F.50 Source 45. Purpose, results, and provenance

Source 45, delivered Section 1.

Added in the eighth write, 7 October 2026. Delivered title block. “Rigidity gaps in Gowers’s Szemerédi framework / Exact additive-energy levels, punctured cosets, and a real fourth-uniformity inequality / Research report prepared with ChatGPT for Vladimir Reshetnikov / 7 October 2026”; *the PDF author field reads the same*. The title page is dated 7 October 2026; the archive arrived at 17:55 on 6 October 2026 (Pacific time). Archive `Gowers_Rigidity_Gaps_Research.zip` (arrival `f70d42385`), manuscript `gowers_rigidity_gaps.tex` with nine section files and a BibTeX database (3,589 lines in ten files, 45 pages), pin `6ce33539e` (its `data/source_manifest.json`, shipped as `data/45-rigidity-gaps-source_manifest.json`). Printed below: the abstract, the status statement after it and Section 1; Sections 2–3 form a chapter of Part II (Section II.183), Sections 4–5 a chapter of Part III (Section III.89), Section 6 and Appendices A–C are in Part V. Its labels carry the stems `en:`, `co:`, `un:`, `col:` and `ap:` of its own sections (*not* this report’s Part prefixes), kept under this report’s prefix. Its ‘question’ environment (“Question”, numbered on its own counter within sections) is set in this report’s ‘Question’ style; its Questions 6.1–6.16 keep their indices. Its BibTeX citations (style `alpha`) point into this report’s bibliography.

Abstract of source 45 (as delivered). We develop three local refinements of the additive-energy and cube-counting framework in Gowers’s proof of Szemerédi’s theorem, with complete proofs and reproducible exact-arithmetic certificates. First, for every $m \geq 5$, the three largest distinct additive energies of m -element subsets of a torsion-free abelian group are M_m , $M_m - 4(m - 2)$, and $M_m - 8(m - 3)$, where $M_m = (2m^3 + m)/3$. We classify every set at the second value using a unit-defect theorem for positive Schur triples. The first gap has a prior derivation from inverse Pollard theory; the additional classification and third gap are the refinements developed here.

Second, writing $\varepsilon = 1 - E(A)/|A|^3$, we extend the existing sharp coset-distance envelope to $\varepsilon < 1/9$ in every abelian group. We identify the canonical closest coset, classify every positive-defect equality case as a coset with a subcoset removed, and prove a second, recursive stability envelope with optimal asymptotic constant one. Its equality cases are described by two nested subcosets. An analogous equality theorem for coordinate-cube collisions forces the deletion into one coordinate.

Third, for real centered functions on finite abelian groups of odd order, we prove $\|f\|_{U^4}^{16} \geq (4795/832)\|f\|_{U^2}^{16}$. An exact example on $\mathbb{Z}/5\mathbb{Z}$ and a dual-cube calculation rule out all real single-character functions as global extremizers of the corresponding norm quotient. The resulting interval is $0.55743364\dots \leq c_4^{\text{odd}} \leq 0.64540701\dots$. We give the resulting improvement to the quartic term in Gowers’s cube bound, a source and normalization crosswalk for ProveIt, and a research agenda. The optimal fourth-uniformity constant and a new global Szemerédi bound are not established.

Status. This is an AI-assisted research manuscript with complete mathematical arguments and supplementary finite checks. It has not undergone external peer review or proof-assistant

verification. Priority claims are limited to the comparisons explicitly documented below. The repository audit is pinned to the commit recorded in Appendix V.169.

Gowers’s proof of Szemerédi’s theorem connects uniformity norms, additive quadruples, structural inverse theorems, and density increments [1]. Several inequalities in that chain have useful extremal questions of their own: what is the first loss from equality, which sets produce it, and which features of an extremizer survive at the next uniformity level? These questions can be answered locally without reworking the entire density-increment iteration. This report studies three of them.

The starting point is the Ramsey directory of ProveIt and its existing research collection [202]. That collection already contains many refinements of the original paper. Reproving a result from an earlier report and presenting it as a new improvement would obscure the actual mathematical progress. We therefore fix a repository revision, inspect the relevant prior reports, and distinguish three kinds of assertions: credited inputs, theorems established in this manuscript, and remaining questions. The proofs of the new assertions do not depend on the finite searches supplied with the article.

F.50.1 The three endpoints

For a finite subset A of an abelian group, its additive energy is

$$E(A) = \#\{(a, b, c, d) \in A^4 : a + b = c + d\}.$$

There are two different cardinality-constrained extremal problems. In a torsion-free group the maximum for $|A| = m$ is

$$M_m = \frac{2m^3 + m}{3},$$

attained by arithmetic progressions. In a general group the universal maximum is m^3 , attained by finite subgroup cosets. The two maxima represent different structures and cannot be interchanged by a change of normalization.

Section II.184 treats the torsion-free maximum. For $m \geq 5$, put

$$H_m = \{0, 1, \dots, m-2, m\}, \quad J_m = \{0, 1, \dots, m\} \setminus \{2\}.$$

We prove that the three largest energy values are precisely

$$M_m, \quad M_m - 4(m-2), \quad M_m - 8(m-3). \quad (\text{F.50.1})$$

At the second value every set is a translate and dilation of H_m , with reflection allowed. The set J_m attains the third value. A full classification at the third value remains a separate problem. The key certificate is the exact integer identity

$$\frac{M_m - E(A)}{4} = \#\{x < y < z \text{ in } A : x + z - y \notin A\}, \quad (\text{F.50.2})$$

first proved for real sets and then transported to arbitrary torsion-free abelian groups. It turns a near-extremal inequality into a classification of a small number of failed additive completions.

Section II.185 treats the coset maximum. For

$$\varepsilon = 1 - \frac{E(A)}{m^3}, \quad \tau(s) = \frac{1 - \sqrt{1 - 4s}}{2},$$

we prove, throughout $0 \leq \varepsilon < 1/9$, that there is a unique closest finite subgroup coset C and

$$\frac{|A \Delta C|}{m} \leq \tau(\varepsilon). \quad (\text{F.50.3})$$

Its subgroup is obtained directly from the autocorrelation: $H = \{d : |A \cap (A + d)| > m/2\}$. Equality with positive defect occurs exactly when $A = C \setminus K$ and K is a subcoset of index at least nine in C . A remainder term measures both points outside C and the failure of the hole K to be a coset. More precisely, set $\delta = |A \Delta C|/m$ and $\eta = \varepsilon - \delta(1 - \delta)$. When $\delta > 0$ and $0 \leq s = \eta/\delta^3 < 1/9$, we prove the second envelope

$$\Delta_C(A) := \frac{1}{m} \min_{D \subsetneq C} |A \Delta (C \setminus D)| \leq \delta \tau(s),$$

where D ranges over nonempty subgroup cosets. Its optimal asymptotic leading constant in the scale η/δ^2 is one. All positive equality cases are $(C \setminus D) \cup J$, with two nested subcosets $J \subsetneq D \subsetneq C$ and inner index at least nine. For coordinate-cube collisions, positive-defect equality in the original coset-distance envelope requires a Cartesian product with just one punctured coordinate.

Section III.90 concerns a functional endpoint. With normalized Gowers norms, define

$$c_d^{\text{odd}} = \sup_{\substack{G \text{ finite abelian of odd order} \\ f: G \rightarrow \mathbb{R}, \mathbb{E}f=0, f \neq 0}} \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4}. \quad (\text{F.50.4})$$

The existing report establishes $c_3^{\text{odd}} = 1/\sqrt{2}$ and asks for the higher-dimensional constants [202]. We prove

$$\left(\frac{39037448}{404304625} \right)^{1/4} \leq c_4^{\text{odd}} \leq \left(\frac{832}{4795} \right)^{1/4}. \quad (\text{F.50.5})$$

The upper estimate comes from a mean-sensitive U^3 inequality, positive Fourier coefficients of an autocorrelation, and a rational quadratic certificate. The lower estimate is attained by the particular test function $(2, 1, -2, -2, 1)$ on $\mathbb{Z}/5\mathbb{Z}$, without any assertion that it maximizes the quotient. A dual-cube identity supplies an improving direction at every real single-character function whose character has odd order at least five; order three is treated separately. Thus the single-character extremal model valid for the preceding norm comparison does not answer the fourth-uniformity problem.

F.50.2 What is inherited and what is added

The baseline and first energy gap are not assigned a new priority claim. The first gap is a short consequence of the equality classification in Pollard's theorem [190]; we include an elementary endpoint proof because it leads efficiently to the second-level classification and third-level exclusion. Hegyvári's energy-spectrum paper already provides the endpoint recurrence, the congruence modulo four, and the construction producing that gap [189]. The second-level classification and third exclusion interval were not found in the primary sources inspected. That statement records the extent of the audit; it does not establish worldwide publication priority.

Shao's work on large additive energy gives structural conclusions in a different, macroscopic regime [191]. Our exact torsion-free window has an absolute deficit of order m , hence relative deficit of order m^{-2} from M_m . It should not be read as a replacement for constant-error inverse theory. Likewise, the extremal-norm results of Eisner and Tao concern near-maximal Gowers norms under magnitude-norm constraints [7]; the real centered quotient in (F.50.4) imposes a different comparison.

Problem	Credited starting point	Extension proved here
Torsion-free energy	The progression maximum in Gowers, Section 7; the first gap follows from inverse Pollard; endpoint calculations also appear in Hegyvári	Complete second-level classification, an exact third energy value, and a unit-Schur-defect theorem
Coset approximation	The sharp quadratic envelope for $\varepsilon \leq 1/100$ in report source 14	The range $\varepsilon < 1/9$, canonical subgroup, all equality cases, and quantitative stability of the hole
Coordinate-cube collisions	The product-coset result in the same small range	The extended range and all positive-defect equality domains
Real uniformity norms	$c_3^{\text{odd}} = 1/\sqrt{2}$ and consequently $c_4^{\text{odd}} \leq 1/\sqrt{2}$ in the report	A smaller upper bound, an exact five-point lower bound, and exclusion of all single-character global extremizers
Cube counting	Exact quartic coefficient P_d and centered expansion in prior report sources	The improved coefficient $P_d(832/4795)^{1/4}$ for $d \geq 4$

Table F.50.1: Scope of the comparisons. The repository reference is [202]; the published energy sources are [1, 190, 189].

F.50.3 Conventions and logical scope

All groups are abelian and written additively, except that characters take values multiplicatively on the unit circle. A coset in a distance minimization is a coset of a finite subgroup. All unmarked cardinality counts are ordered counts, including degenerate additive cubes. On a finite group G we use probability averages $\mathbb{E}_x = |G|^{-1} \sum_{x \in G}$ and normalized Fourier coefficients

$$\widehat{f}(\xi) = \mathbb{E}_x f(x) \overline{\xi(x)}.$$

The dual-group sums are unnormalized. For $d \geq 1$,

$$\|f\|_{U^d}^{2^d} = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{\omega \in \{0,1\}^d} \mathcal{C}^{|\omega|} f(x + \omega \cdot h), \quad (\text{F.50.6})$$

where $\mathcal{C}$ denotes complex conjugation. In the real-valued sections the conjugations may be suppressed. We use the standard cube Cauchy–Schwarz inequality, Fourier orthogonality, and monotonicity $\|f\|_{U^d} \leq \|f\|_{U^{d+1}}$; these are the elementary Gowers-norm facts developed in Section 3 of [1].

The torsion-free, arbitrary-group, and odd-order hypotheses belong to different theorems. In particular, the ordered energy maximum cannot be applied to an arbitrary subset of a cyclic group without rectification, and the real norm comparison cannot be extended to arbitrary complex functions: a nontrivial character has U^j norm equal to one for every $j \geq 2$. All finite searches below are diagnostic or certify a specified finite sum. The universal statements are proved analytically. Their local improvements have not been propagated through Gowers’s complete iteration, so no new numerical upper bound for the global Szemerédi threshold or for $r_k(N)$ is claimed.

[write] **7 October 2026; sources 40–46.** *Identifications and credit within this report.* [202] is this repository at 6ce33539e; at that pin this report’s article held sources 01–14 (“The consolidated article at this pin includes sources 01–14”, Appendix V.169, was true then); “report source 14” and “source 15” are sources 14 and 15 here. *Not cited, and credited here:* sources 28, 29, 34 and 35, whose archives were in docs/incoming/ at the pin. The dual identity $\mathcal{D}_4\phi_1 = 111\phi_1 + 27\phi_3 + \phi_5$, the value 222 and the single-pair non-extremality

(Section III.90) are source 28’s Lemma III.70.3 and Theorem III.70.9 and source 34’s Lemma III.71.3; the even-moment bound is sources 29’s and 34’s $C_m S^m$; the lower endpoint $(39037448/404304625)^{1/4} = 0.5574336$ is weaker than source 28’s and source 34’s 0.5587294 and the batch-121 certificate’s 0.5587725; the mean-sensitive lemma keeps the zero frequency in the three-pairing count of sources 21, 35 and 38, credited only as “the pinned report”. What is new: the coefficient 4795/832, the best among sources 01–46, the mean-sensitive lemma, the energy levels and the coset range. The interval in (F.50.5) is correct but not the best bracket, which is $0.5587725 \leq c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$ (item III.1).

F.51 Source 46. Research target, contribution, and boundaries

Source 46, delivered Section 1.

Added in the eighth write, 7 October 2026. Delivered title block. “Sharp Cubical Correlation and Characteristic Barriers in Gowers’ Polarization Step / Exact binary cubic norms, improved restriction bounds, optimal phase depth, and trace-multiplication tensors / Research contribution prepared for the ProveIt project / 6 October 2026”; *the PDF author field reads “Research contribution prepared for ProveIt”; no statement on AI assistance is made.* Archive `gowers_polarization_barriers.zip` (arrival `f2baec573`), manuscript `article.tex` (1,503 lines, 23 pages); *no pin:* its `SOURCE_AUDIT.md` (shipped as `46-polarization-barriers-SOURCE_AUDIT.md`) records two blob identifiers. Printed below: the abstract, its keywords and Section 1; Sections 2–9 form a chapter of Part IV (Section IV.25), Sections 10–12 and Appendix A are in Part V. Its ‘question’ environment (“Research question”, on the shared counter, in italic type) is set in this report’s ‘Research question’ style, and its remarks, which it sets in upright type, in this report’s remark style.

Abstract of source 46 (as delivered). Let $T : V^3 \rightarrow \mathbb{F}_2$ be symmetric and trilinear, and let $C_T(a, b) = T(a, a, b) + T(a, b, b)$ have rank r . We prove that the exact norm of the T -twisted, eight-function cubic average on the unit ball of $\ell^\infty(V)$ is $2^{-r/2}$. Equality is attained with all eight functions equal to one nonclassical cubic phase taking values among the eighth roots of unity. The proof converts the cube into a family of quadratic convolution operators; a symplectic decomposition then supplies matching extremizers. Consequences include an exact approximate-polarization threshold, a spectral stability identity, and the improvement of the codimension bound in Tidor’s symmetric seven-function restriction proposition from $8 \log_2(1/\delta)$ to $2 \log_2(1/\delta)$.

To connect these results to Gowers’ Lemma 17.1 and Proposition 17.2, we give explicit all-characteristic scalar polarization with the optimal root-of-unity alphabet. For $T_{c,d}(h_1, \dots, h_d) = \text{Tr}(ch_1 \cdots h_d)$ over $\mathbb{F}_{p^n}$, we classify exact nonclassical integrability in every degree, compute the first obstruction’s rank, and determine the largest integrable subspaces and the exact number of elementary symmetric corrections. The general integration criterion is established prior work; the paper uses it with explicit proofs in the degrees needed here. All analytic assertions are proved, rather than inferred from experiments. Reproducible exact finite checks, a source audit, and a formalization roadmap accompany the article. No improved global bound for Szemerédi’s theorem or exhaustive priority claim is made.

Keywords. Gowers norms; polarization; low characteristic; nonclassical polynomials; symmetric multilinear forms; finite-field trace; symplectic restriction; sharp correlation bounds.

The late stages of Gowers’ proof of Szemerédi’s theorem turn structured multilinear frequency information into polynomial phases. The familiar polarization formula divides by a factorial. In the prime-field setting this is harmless when the characteristic is larger than the degree; outside that range it conceals two different phenomena. Sometimes the appropriate repair is merely to enlarge the phase alphabet. Sometimes no circle-valued phase can integrate the given tensor at all.

The corrected source in the ProveIt Ramsey directory explicitly assumes $k < N$ in Lemma 17.1 and $k + 1 < N$ in Proposition 17.2, with N prime [1, 203]. We investigate this boundary rather than changing a constant in an unrelated early density-increment argument. The central question is quantitative: when exact integration fails, how much twisted cubical correlation remains possible, and how much restriction or correction is necessary to remove the obstruction?

F.51.1 Main conclusions

For a symmetric binary trilinear form, a single alternating bilinear form controls several a priori different quantities. Writing $r = \text{rank } C_T$ and $m = r/2$, we prove

$$\sup_{\|f_\varepsilon\|_\infty \leq 1} |\Lambda_T((f_\varepsilon))| = 2^{-m}, \quad (\text{F.51.1})$$

$$\min\{\text{codim } W : T|_{W^3} \text{ integrates}\} = m, \quad (\text{F.51.2})$$

$$\max \Pr(\text{a vertex-separated polarization identity holds}) = \frac{1 + 2^{-m}}{2}. \quad (\text{F.51.3})$$

The average Λ_T is defined precisely in Section IV.27. The first supremum is unchanged if all eight functions are required to be a single suitable eighth-root cubic phase. The second identity follows from alternating-form geometry; the first and third supply exact analytic counterparts, not just a qualitative obstruction.

A related seven-function correlation Γ_T satisfies

$$|\Gamma_T| \leq 2^{-r/4}. \quad (\text{F.51.4})$$

Consequently, correlation at least $\delta > 0$ forces an integrable restriction of codimension at most $2 \log_2(1/\delta)$. This improves the bound $8 \log_2(1/\delta)$ in Proposition 4.6 of Tidor [150], under exactly its symmetry and boundedness hypotheses. We do *not* assert that the new seven-function bound is sharp; the eight-function bound is sharp.

For extension-field trace tensors, we obtain a complete exact-integration classification and a closed formula for m . These give explicit families in which the required codimension grows linearly with the extension degree, together with matching sharp cubical norms. The scalar prime-field case, in contrast, always integrates after passing to an optimal higher-order root-of-unity alphabet.

F.51.2 What is inherited and what is proved here

The theory of nonclassical polynomials is essential background, not a new invention of this note. Tao–Ziegler establish the necessary repeated-variable conditions for total derivatives and the coordinate description of nonclassical polynomials [184]. Tidor proves the converse integration statement for all orders [150, Proposition 3.5]. Quadratic Gauss-sum magnitudes, symplectic decomposition, and Gowers–Cauchy–Schwarz are standard ingredients.

The main results proved in this note are the exact mixed cubic norm and its explicit extremizers, the quantitative agreement and stability consequences, the sharpened symmetric seven-function restriction estimate, and explicit trace-tensor classification and repair formulas. The scalar constructions are effective refinements of the established nonclassical framework. They are included because they identify precisely what an all-characteristic replacement of Gowers’ polarization step must say.

The source audit also records Milićević’s 2026 preprint on general U^4 inverse theory [187]. Its existence makes an unrestricted priority claim especially inappropriate: the present argument is

self-contained, but we have not exhaustively compared every result in that large preprint or the entire literature. “Proved here” describes the content of this manuscript, not a certified first publication. Nor do these local improvements by themselves alter the global bound for progression-free sets. Recent global work such as [2] concerns a different scale of quantitative improvement.

[write] **7 October 2026; sources 40–46.** *Identifications and credit within this report.* [203] is the transcription (blob 41085a0) and [204] this directory’s README (blob fe4d31eb, current from b16ce380d to eae64ceca); source 46 records no commit. *Not cited, and credited here:* source 36 (Section IV.15), the odd-characteristic counterpart, whose archive was in docs/incoming/ while that README was current; and source 43, of the same day, which proves the main theorem (Theorem IV.27.1) and the codimension $2\log_2(1/\delta)$ (Theorem IV.28.1) independently (notes there). [187] is a different preprint of Milićević from [32]. Tidor’s Proposition 4.6, re-read at the write in the arXiv HTML of version 2, displays the codimension $8\log_2(1/\delta)$ quoted here. The intake checked Theorem IV.27.1 exhaustively on $\mathbb{F}_2^2$ (all 16 forms, all 4,096 functions with values among the eighth roots of unity) and on 24 forms over $\mathbb{F}_2^3$ against all 65,536 nonclassical cubic phases; source 46’s suite was rerun on a copy at the write, with identical output.

F.52 Source 47. The problem and the main results

Source 47, delivered Section 1.

Added in the ninth write, 7 October 2026. Delivered title block. “Exact Boolean Obstruction Energies / Derivative periods, sharp quartic–sextic integration, and dimension-free higher-degree extremizers / Research manuscript prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt project / 6 October 2026”; *the PDF author field reads* “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov”. Archive Boolean_Obstruction_Energies_Research.zip (arrival b6a8ee966), manuscript article.tex (1,512 lines, 23 pages), pin 5c9a442d2 (an “observed branch ref”; its SOURCE_AUDIT.md, shipped as 47-boolean-energies-SOURCE_AUDIT.md, records two blob identifiers). It continues source 43 (note in this section). Printed below: the abstract, the status paragraph of its title page and Section 1; Sections 2–9 form a chapter of Part IV (Section IV.42), Sections 10–13 and Appendix A are in Part V. Its macros \F (the field $\mathbb{F}_2$), \Torus ($\mathbb{R}/\mathbb{Z}$), \Sym and \ip are set as this report’s \Ftwo, \Tcirc, \sym and \ipthirtyone, which print the same; its ‘question’ environment (“Research question”, numbered 1–9 on its own counter) is set in this report’s ‘Research question’ style, numbered within its section.

Abstract of source 47 (as delivered). A selected derivative Fourier coefficient carries constraints that disappear when the derivative is replaced by an arbitrary bounded function. Over $\mathbb{F}_2^n$, every even sum of the derivative directions is an exact translation period. Retaining this elementary constraint resolves the higher-dimensional quartic extremal question posed in the preceding Boolean phase-integration draft. More generally, for the symmetric tensor $C_d = \sum_i v(h_i) \prod_{j \neq i} u(h_j)$ with u, v independent, we prove that its maximum selected energy is exactly $1 - (d+1)/2^d$ in every degree $d \geq 3$ and every dimension $n \geq 2$; arbitrary optimizing functions may use all the coordinates. A defect classification and a sharp $7/32$ support theorem for non-pure symmetric trilinear maps then give the exact universal symmetric nonintegrability thresholds $11/16$, $13/16$, and $57/64$ in degrees four, five, and six. In all degrees the universal upper bound is $1 - \min\{d+1, 7\}/2^d$, improving the previous repeated-variable rank-profile bound. We also classify one-dimensional defect images, give a nonnegative deficit recursion, and retain explicit margins for corrupted selectors. Full proofs of the bounds and the required integration and shear identities are included. Nonclassical integration and the earlier shear method are credited as background; the new claims are accompanied by exact finite checks and a formalization plan.

Scope and status. Unrefereed research manuscript. The quartic question is a question in the preceding project draft, not a claim about a previously published famous open problem. Literature

priority and proof-assistant verification are not asserted. These are local phase-integration results, not a new end-to-end bound for Szemerédi's theorem.

F.52.1 The selected energy

Throughout, $V = \mathbb{F}_2^n$ is a finite-dimensional Boolean vector space, all averages are normalized, and

$$\partial_h f(x) = f(x+h)\overline{f(x)}, \quad \Delta_h P(x) = P(x+h) - P(x), \quad e(t) = \exp(2\pi it).$$

For a genuinely d -linear map $T : V^d \rightarrow \mathbb{F}_2$, define

$$\mathcal{E}_T(f) = \mathbb{E}_{h_1, \dots, h_{d-1}} \left| \mathbb{E}_x \partial_{h_1} \cdots \partial_{h_{d-1}} f(x) (-1)^{T(h_1, \dots, h_{d-1}, x)} \right|^2. \quad (\text{F.52.1})$$

The last slot of T is the frequency-evaluation variable. The assumption $\|f\|_\infty \leq 1$ will always be stated or in force; under it, $0 \leq \mathcal{E}_T(f) \leq 1$. Write

$$\mathcal{M}(T) = \max_{\|f\|_\infty \leq 1} \mathcal{E}_T(f).$$

This is a maximum, since a finite product of closed complex unit discs is compact and the energy is continuous.

A symmetric d -linear form is *integrable* if it is the constant top derivative of a degree-at-most- d nonclassical polynomial $P : V \rightarrow \mathbb{R}/\mathbb{Z}$:

$$\Delta_{h_1} \cdots \Delta_{h_d} P(x) = \frac{1}{2} T(h_1, \dots, h_d). \quad (\text{F.52.2})$$

Nonclassical degree means that all $(d+1)$ -fold additive differences vanish. Denote the subspace of integrable tensors by $\mathcal{I}_d(V)$. This is the space usually denoted $\text{nCSM}^d(V)$, not the space of all symmetric multilinear forms. Section IV.43 proves the precise criterion and an explicit primitive.

F.52.2 A dimension-free extremal formula

For independent $u, v \in V^*$, put

$$C_d^{u,v}(h_1, \dots, h_d) = \sum_{i=1}^d v(h_i) \prod_{j \neq i} u(h_j), \quad b_d = 1 - \frac{d+1}{2^d}. \quad (\text{F.52.3})$$

We normally omit the superscripts. The tensor depends on two coordinates; the function in its extremal problem need not.

Theorem F.52.1 (Exact canonical energies). *For every $d \geq 3$, $n \geq 2$, and independent $u, v \in V^*$,*

$$\mathcal{M}(C_d^{u,v}) = b_d = 1 - \frac{d+1}{2^d}. \quad (\text{F.52.4})$$

The constant function attains equality. The same maximum holds for $C_d^{u,v} + S$ whenever $S \in \mathcal{I}_d(V)$.

The proof is in Section IV.46. It does not assume that an optimizer factors through $V/(\ker u \cap \ker v)$, and it does not obtain an upper bound by tensoring a two-dimensional computation.

F.52.3 Sharp low-degree thresholds and a uniform higher-degree bound

Define the universal symmetric nonintegrable extremum in dimension n by

$$M_{d,n}^{\text{sym}} = \max_{T \in \text{SymMult}^d(V) \setminus \mathcal{I}_d(V)} \mathcal{M}(T), \quad n \geq 2. \quad (\text{F.52.5})$$

Here $\text{SymMult}^d(V)$ denotes symmetric multilinear *maps*; no use of symmetric tensor coinvariants in characteristic two is intended.

Theorem F.52.2 (Universal bounds and exact quartic–sextic thresholds). *For every $d \geq 3$ and $n \geq 2$,*

$$\boxed{1 - \frac{d+1}{2^d} \leq M_{d,n}^{\text{sym}} \leq 1 - \frac{\min\{d+1, 7\}}{2^d}.} \quad (\text{F.52.6})$$

In particular, independently of dimension,

$$M_{3,n}^{\text{sym}} = \frac{1}{2}, \quad M_{4,n}^{\text{sym}} = \frac{11}{16}, \quad M_{5,n}^{\text{sym}} = \frac{13}{16}, \quad M_{6,n}^{\text{sym}} = \frac{57}{64}. \quad (\text{F.52.7})$$

If a symmetric tensor has energy strictly above the corresponding upper bound in (F.52.6), it is integrable. For a primitive P_T one then has the lossless identity

$$\mathcal{E}_T(f) = \|fe(-P_T)\|_{U^d}^{2^d}. \quad (\text{F.52.8})$$

The strict thresholds in degrees three through six cannot be decreased.

[write] **7 October 2026; sources 60–69.** *Exact in every degree* (sources 65 and 69, Section IV.51): the maximum over nonintegrable symmetric tensors equals $1 - (d+1)/2^d$ for every $d \geq 3$, and its maximizers are the canonical tensors modulo integrable ones (their Theorems F.70.1 and F.74.1); at $d = 7$ the value is $15/16$.

The distinction between *energy* and the Gowers norm itself is important: the normalization in (F.52.1) corresponds to the 2^d th power of the U^d norm after phase removal.

Degree	Canonical maximum	Previous profile upper bound	New universal upper bound	Exact?
3	1/2	1/2	1/2	yes
4	11/16	3/4	11/16	yes
5	13/16	7/8	13/16	yes
6	57/64	15/16	57/64	yes
7	15/16	31/32	121/128	not established
8	247/256	63/64	249/256	not established

The previous column is the bound $1 - 2^{2-d}$ from the repeated-variable rank profile in the preceding draft [205]. It is a comparison with that specific local estimate, not a claim about the best published global inverse theorem. For $d \geq 7$, the gap left here is exactly $(d-6)/2^d$ between our lower and upper bounds.

F.52.4 What is new relative to the inspected sources

Gowers’s Section 17 turns structured derivative frequencies into a phase removal step [1]. In small characteristic, symmetry alone no longer characterizes top derivatives; the appropriate nonclassical integration theorem is established background [150]. The present problem asks how much selected energy a genuinely nonintegrable symmetric tensor can support.

The preceding project manuscript [205] proved the cubic rank law, the general shear profile, the two-dimensional quartic value $11/16$, and a reduction of the universal quartic problem to the pullback of C_4 . Its Research Question 1 asked whether the $11/16$ value persists when the optimizing function may depend on extra coordinates. Theorem F.52.1 answers that question affirmatively and gives an all-degree formula. Theorem F.52.2 extends the sharp universal conclusion through degree six. The derivative-constrained rank-one estimate, the pure-defect classification, the support argument, and the ensuing deficit recursion are the new proof components developed here. Background components are reproved so that no unreviewed earlier theorem must be accepted on trust.

[write] **7 October 2026; sources 47–53.** *Identifications within this report.* [205] is source 43 (Section F.48), read from the user’s library under its file name, not by repository path; its “Research Question 1” is source 43’s Research question V.160.1, and its “previous profile” bound $1 - 2^{2-d}$ is source 43’s repeated-variable rank profile. [206] is this repository; the observed ref 5c9a442d2 is an ancestor of HEAD, and the two blobs recorded in 47-boolean-energies-SOURCE_AUDIT.md (Sections17_18.lean 18495f5, Definitions.lean 97113b7) are unchanged at HEAD 06b30ca52. *Source 43’s quartic question is answered:* Theorem F.52.1 at $d = 4$ gives $\mathcal{M}(C_{4,n}) = 11/16$ for every $n \geq 2$, where source 43 had proved $11/16 \leq M_{4,n} \leq 3/4$ (item IV.10). The thresholds of Theorem F.52.2 concern *symmetric* tensors (the evaluation slot included); for general d -linear forms over $\mathbb{F}_2$ the sharp certification threshold is source 43’s $1 - 2^{1-d}$ (Theorem IV.38.4), which is larger in every degree $d \geq 3$, as it must be. Checked at the write in exact rationals: $\mathcal{E}_{C_d}(1) = 1 - (d+1)/2^d$ for $d = 3, \dots, 10$ by direct enumeration on $\mathbb{F}_2^2$, and every entry of the table above ($b_7 = 15/16 = 120/128$, $b_8 = 247/256$; the upper bounds $121/128$, $249/256$; the previous column $1 - 2^{2-d}$); the intake checked the canonical maximum and the thresholds through degree six by random-restart local search and exhaustive form lists (evidence, not proof). Source 47’s suite was rerun on a copy at the write, normally and with `-O`, with identical output.

F.53 Source 48. The research target and the resulting strengthening

Source 48, delivered Section 1.

Added in the ninth write, 7 October 2026. Delivered title block. “Collision-aware derandomization of Gowers’s restriction step / Exact energy inventories, prescribed-cardinality subsets, and constructive restriction for linear systems / Research manuscript prepared with ChatGPT for the ProveIt project / 6 October 2026”; *the PDF author field reads the same.* Archive `Gowers_Collision_Aware_Derandomization.zip` (arrival b6a8ee966), manuscript `article.tex` (1,763 lines, 27 pages), pin 0b4890a1c (its `SOURCE_AUDIT.md`, shipped as `48-derandomization-SOURCE_AUDIT.md`). It continues source 17 (note in this section). Printed below: the abstract, its keywords and Section 1; Sections 2–9 form a chapter of Part II (Section II.186), Sections 10–13 and Appendices A–B are in Part V. Its macros `\cost` (C_s , not this report’s calligraphic $\mathcal{C}$) and `\pin` are renamed `\costfortyeight` and `\pinfortyeight`; its ‘question’ environment (“Research question”, on the shared counter) is set in this report’s ‘Research question’ style.

Abstract of source 48 (as delivered). Gowers’s random-restriction argument improves the proportion of additive relations respected by a function. Turning its existence proof into an explicit algorithm requires conditional expectations of high-order tuple counts. Ordinary

weighted convolution is not sufficient: a repeated label is selected once, not independently at each occurrence.

We construct an exact finite group-algebra inventory that records these collisions without enumerating tuples. For order s , n labels and ambient group size D , it evaluates the Bernoulli expectation and deterministically rounds an arbitrary rational signed combination of energy counts in $O(s^4 n D)$ rational operations and $O(s^2 D)$ storage. A single branching decision needs only $O(s^2)$ coefficient probes; the remaining work consists of explicit translations of coefficient arrays. A distinct-support extension gives the corresponding theorem for a uniformly sampled subset of any prescribed cardinality, with no multiplicative dependence on that cardinality. The method extends to arbitrary fixed systems of linear equations, including zero coefficients, repeated entries, and composite moduli; repeated column types substantially reduce its complexity.

As a source-specific application, we give a fully deterministic, rational-arithmetic version of polynomial-loss Fejér restriction. It retains the existing bound $|B|^{2s-1}/(2\eta L^{2s})$ with no additional loss. The complete prime-cyclic search takes $O_s(np^5)$ rational operations, whereas rounding after a filter is fixed takes $O_s(np^2)$. These two costs are distinguished throughout. The contribution resolves a specific algorithmic question in the inspected companion research, not a new global bound for Szemerédi’s theorem. Complete proofs, exact tests through order eight, executable reference code, and integration notes are included. Literature priority and formal Lean verification are not claimed.

Keywords. Additive energy; Gowers restriction; conditional expectation; repeated entries; divided powers; fixed-size sampling; constructive additive combinatorics.

F.53.1 What is being made constructive

In Lemma 9.3 of his 2001 paper, Gowers restricts a function $\phi : B \rightarrow \mathbb{Z}/p\mathbb{Z}$ so that a large proportion of its surviving additive sixteen-tuples are respected [1]. The random choice is structured: first choose phase parameters, then include distinct points independently with the resulting probabilities. The second stage is independent across points, but not across the occurrences of one point in a tuple.

The inspected companion manuscript [207] replaces the long product of elementary filters by one normalized Fejér kernel and obtains polynomial losses in the initial respected fraction and the desired error. Its further-research question on efficient certification and extraction asks for an algorithm that handles conditional Bernoulli choices, repeated entries, and vector-valued graph data without listing all tuples. Its existence argument explicitly does not supply such a running-time theorem. We address that question for balanced additive relations, and then for arbitrary fixed linear systems.

The method of conditional expectations is classical [208]. Our contribution is not the averaging principle itself. It is the exact coefficient representation, the local replacement identity, the resulting complexity bound, the prescribed-cardinality extension, and the source-specific rational implementation of the full restriction step.

F.53.2 The main computational theorem

For a finite abelian group K and labelled group elements $g_1, \dots, g_n$, put

$$E_s(S; g) = \# \left\{ (i_1, \dots, i_{2s}) \in S^{2s} : \sum_{j=1}^s g_{i_j} = \sum_{j=s+1}^{2s} g_{i_j} \right\}. \quad (\text{F.53.1})$$

Labels, rather than group images, are independently selected. The images g_i need not be distinct. In the graph application the labels are the distinct elements of B , with two images x and $(x, \phi(x))$.

Suppose a statistic is a rational signed sum

$$F(S) = \sum_{\nu=1}^v c_\nu E_s(S; g^{(\nu)}), \quad c_\nu \in \mathbb{Q}, \quad D = \sum_{\nu=1}^v |K_\nu|.$$

The common label set is essential: all energies are evaluated on the same selected subset. Write

$$C_s = \binom{s+2}{2}^2. \quad (\text{F.53.2})$$

Theorem F.53.1 (Headline rounding theorem). *Let each label i initially be selected independently with rational probability $p_i \in [0, 1]$. There is a deterministic exact algorithm returning $S \subseteq [n]$ such that*

$$F(S) \geq \mathbb{E}F(S_p).$$

Using explicit arrays indexed by the groups, its arithmetic cost is $O(C_s n D)$ and its storage is $O((s+1)^2 D)$ rational numbers. After initialization, choosing the value of one indicator requires $O((s+1)^2 v)$ coefficient probes and rational operations; maintaining the full state costs $O(C_s D)$ per changed indicator.

For every prescribed $k \in \{0, \dots, n\}$, there is also a deterministic algorithm returning $|S| = k$ with

$$F(S) \geq \frac{1}{\binom{n}{k}} \sum_{|A|=k} F(A).$$

Its arithmetic cost is $O(C_s(2s+1)^2 n D)$ and its storage is $O((s+1)^2(2s+1)D)$. In particular, there is no factor k in this bound.

The first assertion is proved in Sections II.188–II.190; the second in Section II.191. The bounds are not claims of polynomial time in $\log D$. A group specified by a short modulus can be much larger than its bit description. Dense ambient-group storage is an explicit, and sometimes expensive, part of the algorithm.

F.53.3 The source-specific quantitative conclusion

Let $m = 2s \geq 4$, p be prime, $B \subseteq \mathbb{F}_p$ have size $n = \beta p > 0$, and $\phi : B \rightarrow \mathbb{F}_p$. Write $\mathcal{T}_s(B)$ for all balanced m -tuples in B and $\mathcal{R}_s(B, \phi)$ for those whose images satisfy the same relation. Suppose

$$\mathcal{R}_s(B, \phi) \geq \alpha n^{m-1}, \quad 0 < \alpha, \eta \leq 1, \quad L = \left\lceil \frac{m+1}{\alpha\eta} \right\rceil.$$

Under the explicit size hypothesis

$$p > 2L - 2, \quad p \geq 2L^m(2L - 1)^{m-1}\beta^{-(m-1)}, \quad (\text{F.53.3})$$

Section II.194 constructs, rather than merely proves the existence of, a set B' satisfying

$$\mathcal{R}_s(B', \phi) \geq (1 - \eta)\mathcal{T}_s(B'), \quad (\text{F.53.4})$$

$$\mathcal{T}_s(B') \geq \frac{n^{m-1}}{2\eta L^m} \geq \frac{\alpha^m \eta^{m-1}}{2(m+2)^m} n^{m-1}. \quad (\text{F.53.5})$$

The retention estimate is inherited from the companion Fejér argument; we reprove the needed estimate and preserve it exactly during rational approximation and derandomization. This distinction matters: the new result is a constructive strengthening, not a second claim to the same quantitative improvement.

Task in the prime-cyclic graph setting	Arithmetic cost for fixed s
Evaluate one exact Bernoulli expectation	$O_s(np^2)$
Round after phase parameters are fixed	$O_s(np^2)$
Search all translated Fejér parameters and round	$O_s(np^5)$, hence $O_s(p^6)$
Search unshifted parameters, with a larger size threshold	$O_s(np^4)$, hence $O_s(p^5)$
Round to a prescribed size from uniform sampling	$O_s(np^2)$

F.53.4 Status, novelty, and limitations

The original restriction mechanism, the companion’s polynomial retention bound, and conditional-expectation rounding are prior ingredients. The exact array identities and their algorithmic consequences are proved below. No exhaustive bibliographic priority claim is made for these identities, which are related to familiar exponential generating functions and multilinearization. The manuscript resolves the stated algorithmic subproblem with an ambient-size polynomial bound; it does not resolve near-linear extraction, sparse high-dimensional extraction, or all of the companion’s algorithmic questions.

No new global estimate for progression-free sets is asserted. For example, the modern bound of Leng, Sah and Sawhney uses quasipolynomial inverse theorems and a density-increment strategy [2]; improving the implementation of one older restriction step is not by itself an improvement of that global result. The present manuscript also leaves the separate Section 16 closing-parameter problem untouched.

[write] **7 October 2026; sources 47–53.** *Identifications and credit within this report.* [207] is source 17 (Section F.22), read from the user’s library under its file name: its “Fejér restriction theorem” is source 17’s Theorem II.61.2, whose retention bound $T_m(B') \geq b^{m-1}/(2\eta L^m) \geq \alpha^m \eta^{m-1} b^{m-1}/(2(m+2)^m)$ (II.61.6) is the bound (F.53.5) here, in the special case of an unweighted set $B \subseteq \mathbb{F}_p$ and a scalar map (source 17 allows weights and vector-valued maps on $\mathbb{F}_p$ -vector spaces); its “further-research item 9” is source 17’s Question 9 (Section V.71; item II.23). [209] is this repository at 0b4890a1c; `Proofs09Restriction.lean` has the same blob at HEAD 06b30ca52. *Source 17’s Question 9 is answered in its ambient-size form:* the signed objective is evaluated and rounded deterministically in $O(C_s n D)$ rational operations without enumerating tuples (Theorem F.53.1), and the restriction of Lemma 9.3 is constructed with source 17’s retention unchanged (Theorem II.194.1); a near-linear or $\log D$ -polynomial version is not claimed, as source 48 says. Source 48 treats Lemma 9.3 only: the eighth write’s description of it, from its placement, as “a constructive form of the Fejér selections of sources 17 and 40” overstates its scope, since source 40’s selection at Lemma 15.5 (arrangements, Section II.164) is not derandomized here; its general-system theorem (Section II.192) counts solutions of a fixed linear system and is not applied to arrangements. The intake rechecked the Stirling response identity, the hypergeometric survival law and both delivered examples by enumeration; source 48’s suite and its four example runs were rerun on copies at the write (identical output apart from the elapsed-time field).

F.54 Source 49. Scope, baseline, and principal contributions

Source 49, delivered Section 1.

Added in the ninth write, 7 October 2026. Delivered title block. “Second-Order Degeneracy, Dual Uniformity Bounds, and an Acyclic Model Theorem / New refinements in Gowers’ Szemerédi framework / Research manuscript prepared for Vladimir Reshetnikov with ChatGPT / For integration into the ProveIt Ramsey development / 6 October 2026”; *the PDF author field reads “Research manuscript prepared for Vladimir Reshetnikov with ChatGPT”*. Archive `gowers_rank_dual_refinements.zip` (arrival `b6a8ee966`), manuscript `gowers_rank_dual_refinements.tex` (3,329 lines, 45 pages), pin `8c12c906e` (its `source_1` `edger.json`, shipped as `data/49-rank-dual-source_ledger.json`). Printed below: the abstract, the status statement of its title page and Section 1; Sections 2–6 form a chapter of Part II (Section II.220), Section 7 a chapter of Part III (Section III.92), Section 8 stands in the Part II chapter on Lemma 16.10 after source 50 (Section II.226), Sections 9–11 are in Part V. It uses `cleveref` (options `nameinlink`, `noabbrev`); its references are written out as `cleveref` prints them (lowercase names for `\cref`, capitalized for `\Cref`). Its plot is drawn with `pgfplots`, now loaded by this report. Its ‘question’ environment (“Question”, numbered within sections) is set in this report’s ‘Question’ style.

Abstract of source 49 (as delivered). We develop three refinements of the framework underlying Gowers’ quantitative proof of Szemerédi’s theorem. First, in sufficiently large prime characteristic we classify the derivative-rank-two multiaffine polynomials that are ternary-valued on a Boolean cube. Their translation certificates contribute no new component outside the signed-coordinate arrangement. This determines the second coefficient in the degeneracy probability of two congruent coordinate cubes in every fixed dimension, gives a finite cubic-order remainder bound, and proves exact formulas in dimensions one and two. An exhaustive integer certificate extends exactness to dimension three for every odd prime. The exact low-dimensional formulas extend to every finite field of characteristic at least five. A full-subfield alphabet theorem gives an exact all-dimensional formula in characteristic three. More generally, the prime-field degeneracy probability is eventually a fixed Laurent polynomial in the prime.

Second, the complete fourth dual function of a two-harmonic test improves the real, centered uniformity comparison on every finite abelian group of odd order to $\|f\|_{U^4}^{16} \geq (29/4)\|f\|_{U^2}^{16}$ and $\|f\|_{U^5}^4 \geq (50/29)\|f\|_{U^2}^4$. The fourth factor improves to 15/2 when a largest Fourier coefficient has character order different from five. The finite coefficient lemma is proved by a reproducible exact-arithmetic enumeration.

Third, from seven explicitly stated earlier structural interfaces we prove a replacement cover theorem with recurrences that strictly decrease the active dimension. It separates the global core, graph count, width exponent, and minimum-width threshold, and yields a conditional recursive replacement for the Section 16 multilinear inverse conclusion.

The manuscript distinguishes prior project results, new deductive proofs, computer-assisted finite proofs, and imported hypotheses. It supplies complete certificate generators and a research agenda. No sharp norm constant or new global Szemerédi bound is claimed.

Status. Research manuscript; not peer reviewed or formally verified. The exact finite verifiers are included. The Section 16 replacement is an implication from the stated interfaces, whose deep source proofs are not newly certified here.

Gowers’ proof [1] organizes a quantitative density-increment argument around uniformity norms, additive structure, multiaffine phase models, and the counting of coordinate-cube arrangements. These are separate quantitative interfaces. Improving one of them is a mathematical result, but a stronger final bound for Szemerédi’s theorem requires a complete propagation through the remaining interfaces. This article concentrates on three places where a precise improvement can be proved and independently checked.

The immediate baseline is the ProveIt Ramsey development [210], including its corrected transcription of the paper and the earlier project manuscripts. In particular, the degree-filtration report [212] already identifies the first-order two-cube degeneracy coefficient and the affine translation

arrangement. The prime-torsion report [213] already establishes a balanced Fourier-moment profile and the upper bound $c_4^{\text{odd}} \leq 2/3$. The interpolation report [215] already improves sampling and shared refinement and audits the closing numerical comparison in Section 16. Those ingredients are credited and restated where needed; they are not presented as new discoveries.

F.54.1 The degeneracy result

Let $\delta_{k,1}(p)$ be the proportion of degenerate ordered pairs of congruent k -dimensional coordinate cubes over $\mathbb{F}_p$, in the precise one-arrangement convention of section II.221. Write

$$L_k(p) = 1 - (1 - p^{-1})^k \prod_{i=1}^k \left(1 - \frac{2i+1}{p}\right). \quad (\text{F.54.1})$$

The prior argument gave $\delta_{k,1}(p) = L_k(p) + O_k(p^{-2})$. We replace its possible codimension-two error by a codimension-three error. In particular,

$$\delta_{k,1}(p) = \frac{k^2 + 3k}{p} - \frac{k(3k^3 + 14k^2 + 15k - 14)}{6p^2} + O_k(p^{-3}). \quad (\text{F.54.2})$$

An explicit error constant is given in theorem II.222.4. More strongly,

$$\delta_{k,1}(p) = L_k(p) \quad (k = 1, 2, 3; p \text{ any odd prime}). \quad (\text{F.54.3})$$

The cases $k \leq 2$ have elementary proofs. The case $k = 3$ uses an exact finite certificate, together with an algebraic argument reducing all large primes to that certificate. Thus (F.54.3) is not an extrapolation from checking a collection of primes.

The field extension is explicit. For $q = p^m$ with $p \geq 5$, the same product formula holds at $k \leq 3$ with q in place of p . In characteristic three there is a different, exact formula in every dimension:

$$\delta_{k,1}(3^m) = 1 - (1 - 3^{-m})^k \prod_{i=1}^k (1 - 3^{i-m}). \quad (\text{F.54.4})$$

It follows from a linear-independence count over the prime subfield, not from the rank-two classification; see corollary II.225.5.

The underlying structural point is a classification of derivative rank two in sufficiently large prime characteristic. Total degree is at most derivative rank for a multiaffine polynomial. A polynomial of rank two is therefore quadratic, and its quadratic part splits across two disjoint sets of coordinates. Ternary Boolean values force both sets to have size at most two. The resulting very small list has no translation locus outside the already known signed-coordinate hyperplanes. This answers the earlier exactness question through dimension three and determines the next coefficient in all fixed dimensions.

F.54.2 The real uniformity bounds

All norms in this subsection use normalized counting measure. For a complex function on a finite abelian group, our convention is

$$\|f\|_{U^d}^{2^d} = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{\omega \in \{0,1\}^d} \mathcal{C}^{|\omega|} f(x + \omega \cdot h), \quad (\text{F.54.5})$$

where $\mathcal{C}$ is conjugation. On real functions all conjugations disappear. Define

$$c_d^{\text{odd}} = \sup \left\{ \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4} : |G| \text{ odd}, f : G \rightarrow \mathbb{R}, f \neq 0, \mathbb{E}f = 0 \right\}.$$

The new estimates are

$$c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183098\dots, \quad c_5^{\text{odd}} \leq \frac{29}{50} = 0.58. \quad (\text{F.54.6})$$

For functions whose largest Fourier coefficient can be chosen at a character of order different from five, the fourth bound improves to $(2/15)^{1/4} = 0.6042750794\dots$. The prior upper bound $2/3$ and the new one refer to precisely the same normalized ratio.

The improvement uses a two-harmonic function as a dual test for an *arbitrary* f . The earlier report had used such a function as an example giving a lower bound on the extremal constant. Computing its complete dual reveals a much smaller unwanted Fourier tail and turns that example into a useful upper-bound certificate. A moment bound handles spread-out Fourier mass, while the dual test handles concentrated mass. Character orders five, seven, nine, and eleven are checked exactly rather than excluded. The fifth-order result combines the new dual test with its own higher autocorrelation moment.

The earlier lower example remains valid:

$$\frac{4804}{5465198316262956^{1/4}} = 0.5587293909\dots \leq c_4^{\text{odd}} \leq 0.6094183098\dots \quad (\text{F.54.7})$$

This is an improved interval, not a determination of the exact constant. The lower endpoint is from [213].

F.54.3 The replacement model theorem

The third contribution is a complete new induction from stated source interfaces. Its logical form is

$$\text{I1--I7} \implies \text{a cover theorem with explicit decreasing-dimension recurrences.}$$

The global exceptional set is chosen before the local box and its error tolerance. A separate ambient dimension is retained throughout the induction, so a common-difference box partition is never lifted from fewer coordinates without proof. Four controls remain distinct: the prime threshold N , graph count Q , width exponent c , and minimum width T . For fixed active and ambient dimensions, and $u = a\gamma\varepsilon$, they satisfy bounds of the form

$$\begin{aligned} N &\leq (C/(a\gamma))^C, & Q &\leq (C/u)^C, \\ c &\geq \exp(-(C/u)^C), & T &\leq \exp(\exp((C/u)^C)). \end{aligned} \quad (\text{F.54.8})$$

The exact recurrences and their proof are in section II.235. This conditional result supplies the local multilinear-model conclusion needed at the end of Section 16, with the new controls. The article does not identify those controls with the original printed ones, certify all earlier deep lemmas, or propagate a global progression bound.

Statement	Contribution here	Proof status
Translation stabilizer and affine hyperplanes	Restated baseline	Prior project result; proof included
Rank-two classification and census	New structural refinement	Deductive proof and small explicit table
All- k second coefficient; exact $k \leq 2$	New quantitative consequence	Deductive proof
Exact $k = 3$ at every odd prime	Resolves the next dimension	Exact computer-assisted finite proof
Finite-field extensions; all- k characteristic-three count	New generalization	Deductive proofs from subfield structure
Eventual Laurent polynomiality	New structural application	Deductive proof; finite-field arrangement method credited
$c_4 \leq (4/29)^{1/4}$, $c_5 \leq 29/50$	Stronger project upper bounds	Deductive argument with exact finite coefficient lemma
Balanced moment and anchor sampling	Restated baseline	Prior project results; proofs included
Acyclic cover and inverse implication	New complete recurrence	Proved from explicit I1–I7 hypotheses
Sharp norm constants; all- k exactness; global Szemerédi bound	Further research	No claimed solution

Table F.54.1: What is new and what the proofs establish. “New” is relative to the inspected project sources; it is not an assertion of an exhaustive literature priority search.

F.54.4 Contribution and proof-status ledger

Sources and reproducibility. The repository baseline is the following pinned commit:

8c12c906e4ab5a00001a65de3c882f935a15661a

The corrected source transcription, the integrated earlier article, and the cited supplemental project manuscripts were inspected before selecting these problems. The code accompanying this article generates the finite certificates from definitions; it does not merely check a stored answer. Exact verifiers use integer and rational arithmetic from the Python standard library. Section V.187 describes their role and the intended integration boundary.

Conventions. The letter p denotes a prime in the arrangement and model sections. In the norm section, $p \in (0, 1]$ is explicitly reintroduced as a Fourier mass parameter; group order there need not be prime. A multiaffine polynomial is separately affine in each variable. Binomial coefficients with lower argument larger than the nonnegative upper argument are zero. All asymptotic constants may depend on a fixed dimension; no assertion uniform in a dimension growing with p is implicit. We write $X_A = \prod_{i \in A} X_i$ for a squarefree monomial.

[write] **7 October 2026; sources 47–53.** *Identifications and credit within this report.* The “earlier project manuscripts” are sources of this report, read from the user’s library and identified in `data/49-rank-dual-source_ledger.json` by SHA-256 (all four match the delivered manuscripts; the intake verified them): [212] is source 22 (its “exceptional two-cube case” is Theorem II.85.1), [213] is source 34 (Section F.39), [214] is source 35 and [215] is source 19 (Section F.24); [211] is this report at its third write (sources 01–14, `b16ce380d`), and [210] this repository at `8c12c906e`. Two of the six SHA-256 values in the ledger (the report article and the transcription) hash the file followed by one line feed, as for sources 26 and 32 (recorded, not edited). *Not cited:* source 45, whose archive arrived at 17:55 on 6 October 2026, after this source’s pin (17:11); so its “prior

upper bound $2/3$ ” (source 34) is right at its pin, while among sources 01–48 the best upper bound before this source was source 45’s $(832/4795)^{1/4} = 0.6454070$ (Section III.89). The bound $c_4^{\text{odd}} \leq (4/29)^{1/4}$ (F.54.6) supersedes it, and the bracket among sources 01–53 is $0.5587725 \leq c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$, the lower end being the batch-121 intake certificate on $\mathbb{Z}/7\mathbb{Z}$ rather than the 0.5587294 of sources 28 and 34 quoted in (F.54.7) (note after Theorem III.93.1). The “signed-coordinate arrangement” and the first-order coefficient are source 22’s; the balanced-moment profile is source 34’s (and source 29’s $C_m S^m$); the anchor sampling and the audit of the closing comparison are source 19’s (notes in Sections II.221, III.93 and II.235). Source 49’s three exact verifiers were rerun on copies at the write, with identical certificates (the dimension-three certificate has the SHA-256 it records); the optional NumPy cross-check was not run.

F.55 Source 50. The obstruction and its significance

Source 50, delivered Section 1.

Added in the ninth write, 7 October 2026. Delivered title page. “Combinatorics / Ramsey / Gowers–Szemerédi / Sharp Branch-Cover Profiles and a Counterexample to the Two-Input Form of Lemma 16.10 / Local quantitative obstructions in Gowers’s Szemerédi framework / Research manuscript prepared for the ProveIt project / 6 October 2026”; *the PDF author field reads “Research manuscript prepared for the ProveIt project”; no statement on AI assistance is made.* Archive `ramsey_branch_cover_obstructions.zip` (arrival `b6a8ee966`), manuscript `article.tex` (1,734 lines, 27 pages), pin `588552b79` (its `SOURCE_AUDIT.md`, shipped as `50-branch-cover-SOURCE_AUDIT.md`, records four blob identifiers). Printed below: the abstract, the status statement of its title page and Section 1; Sections 2–9 form, with source 49’s Section 8 after them, a chapter of Part II on Lemma 16.10 (Section II.226), Sections 10–13 and Appendices A–C are in Part V. Its macro `\code{\texttt{\detokenize{#1}}}` is set as this report’s `\filename`, the same definition, and its `\pin` renamed `\pinifty`; its ‘question’ environment (“Research question”, numbered 1–10 on its own counter) is set in this report’s ‘Research question’ style, numbered within its section; its listings are set in this report’s listing style.

Abstract of source 50 (as delivered). We study the minimum number of polynomial graphs required to cover a prescribed fraction of a finite-alphabet function on long arithmetic progressions and on partitions into product boxes. For a word whose symbol masses on a progression are $b_1 \geq \dots \geq b_R$, a list of q degree-at-most- d polynomials covers at most $\max_s \{b_1 + \dots + b_s + (q-s)dR\}$ points. If $b_q \geq dR$, the optimum is exactly $b_1 + \dots + b_q$. Strict inequality gives rigidity and a quantitative deficit for every list containing nonconstant polynomials. We extend this law to product boxes, prove a global polynomial stability bound, and construct words with any prescribed branch-mass profile simultaneously on all long modular progressions. The construction is derandomizable by exact conditional expectations. A second construction improves the sufficient scale to $O(R(d+1)\log p)$ for growing alphabets at fixed relative error.

The main application is a counterexample to the precise two-premise `lemma_16_10` recorded in ProveIt’s pinned Section 16 catalogue. At $k = \theta = \gamma = 1$, and for every sufficiently large prime modulus, the two premises hold for a full-domain function $F(h, y) = a(y)$, but the conclusion fails. The failure persists after arbitrary adaptive partitions into proper boxes and is separated from the required retained mass by a fixed positive margin. This is stronger than the previously documented failure of a numerical comparison. The conclusion concerns the stated abstraction, not Szemerédi’s theorem or all the additional hypotheses in Gowers’s upstream construction. We give a parameter-preserving repair interface, exact finite verification artifacts, a formalization plan, and further research questions. We also prove that the full-domain product property at $\gamma = 1$ forces multiaffinity, excluding these witnesses from that stronger upstream hypothesis.

Status. The results are proved in the mathematical manuscript. The new counterexample is not a Lean-kernel-checked theorem and has not undergone independent review. The elementary probabilistic and root-counting ingredients are classical; no comprehensive priority claim is made for their combination. No improved global bound for progression-free sets is asserted.

F.55.1 Which problem is being resolved?

Gowers’s proof of Szemerédi’s theorem passes through multiple multilinearity and a local-to-global lifting step in Section 16 [1]. In the ProveIt implementation, the catalogue entry `lemma_16_10` isolates two inputs: multiple multilinearity of the cross-sections obtained by fixing the final coordinate, and the fibrewise affine covers supplied by the preceding lemma [218]. The repository already warns that the advertised closing numerical comparison is not justified. A separate file proves an obstruction to that comparison, while explicitly distinguishing it from a counterexample to the assertion [220].

Our principal result closes that particular logical gap in the negative: *the two-input implication itself is false with the recorded quantitative controls*. It is not enough to replace the unsuccessful numerical calculation by a more efficient argument. At least one hypothesis, conclusion, or control function must change.

This is a local research result, rather than a new global Szemerédi bound. Its significance is that the false abstraction sits at a proof interface: trying to prove it, or using it as though it were established, would obstruct a sound completion of the argument. The affirmative results identify the quantitative information that is lost at that interface.

F.55.2 The model example

Let $a : \mathbb{F}_p \rightarrow \{v_1, \dots, v_R\}$ be a word over R distinct field elements and put

$$F(h, y) = a(y), \quad h \in \mathbb{F}_p^k, \quad y \in \mathbb{F}_p. \quad (\text{F.55.1})$$

Every fixed- y cross-section is constant. Every final-coordinate fibre is covered, with no exceptional points, by the R constant affine maps $y \mapsto v_j$. These are exceptionally strong qualitative hypotheses. Nevertheless, a nonconstant affine function takes any one value v_j at most once. On a long progression on which all symbols have nearly equal masses, one constant graph covers about $1/R$ of the word, and a nonconstant affine graph covers at most R points. Thus q graphs cannot cover much more than q/R of the progression. The same statement holds on every cell of a box partition.

The issue is not an inability to find suitable anchors or to interpolate. All R branches are already global constants. The obstruction is the attempted removal of branch complexity from the output budget.

F.55.3 An overview of the results

The exact one-dimensional cover law is Theorem II.228.2; its product-box and stability versions are Theorems II.229.1 and II.229.4. Theorem II.230.2 constructs words balanced on all sufficiently long modular progressions, and Theorem II.230.3 gives an exact deterministic construction. Theorem II.231.1 realizes every finite probability profile as a local polynomial-cover profile.

Theorem II.233.1 then refutes the pinned two-input form of Lemma 16.10, with every witness and numerical inequality written out. A particularly concrete choice is

$$E = 2^{1024}, \quad Q = 2^E, \quad R = 4Q, \quad c = Q^{-1}. \quad (\text{F.55.2})$$

For every prime $p > 2^{4096Q^4}$ there is a word a for which both inputs hold, whereas every permitted output cover retains at most $5p^2/16$, less than the required $p^2/2$. The size is enormous because we instantiate the repository’s own controls. The mechanism has no enormous-constant dependence: the general cover laws apply to small parameters as well.

F.55.4 Scope, sources, and novelty

All source-dependent assertions about the formal catalogue refer to commit

588552b79497c93d0e2f3a3f9224fd7681e8d396

The Section 16 file has blob `c5c7d2bee91bc589de1d3082429f7ee3597e99f9`. A later read of the default branch during this investigation returned the same blob. The article does not assume that an evolving branch remains unchanged.

The new contribution is the branch-cover obstruction, its exact profile and stability formulations, and its fully quantified application to that catalogue entry. Concentration, polynomial root bounds, grid interpolation, and conditional expectation are standard tools. We reprove the versions needed here and do not present those tools themselves as discoveries. The repository already contains positive interpolation and local lifting machinery; we retain its explicit parameters instead of rebranding it as new [222].

[write] **7 October 2026; sources 47–53.** *Identifications and credit within this report.* [218], [219], [220], [221] and [222] are files of this repository at the pin 588552b79; the four recorded blobs match, and `Section16.lean`, `Definitions.lean` and `Sections14_15.lean` are unchanged at HEAD 06b30ca52, while `Proofs16LocalAffineLift.lean` has changed (note in Section II.234). The repository’s “warning” about the closing comparison is source 19’s audit (Section V.80.2) and its formal companion `section16_same_dimension_closing_comparison_fails`. Source 50 names no earlier report; the Section 16 material of this report it bears on is source 19’s sampling (Section II.65), source 23’s gluing (Section II.81) and source 39’s shared anchors (Section II.151), and source 49’s conditional acyclic replacement of the same day follows in this chapter.

Formalized since the pin (checked at HEAD 06b30ca52, ledger 113 exact companions and 7 open statements). The repository has formalized this source’s counterexample along its route (`FORMALIZATION_STATUS.txt` names source 50): the finite-alphabet cover bound and its 5/16 box-partition consequence, balanced words on every long proper progression by a Hoeffding bound and a union bound, the identity-partition certificates for both packaged premises, and the literal negation $\neg \text{lemma_16_10}$, the theorem `lemma_16_10_packaged_premises_counterexample` (`Proofs16PackagedPremisesCounterexample.lean`, line 13; commit c0a404a59, 7 October 2026, 00:47), with $k = \theta = \gamma = 1$, full domains, all cubes and zero anchor. The formal witness takes an existential large prime and a word balanced on progressions of length at least N^δ , not this source’s explicit bound $p > 2^{4096Q^4}$. Beside it the repository keeps a positive interface, `Section16AllBoxLineCovers.explicit_multilinear_cover` (`Proofs16PackagedExplicitCover.lean`, line 11; 1a8581c65), and it has formalized this source’s endpoint characterization (note after Proposition II.234.1). The ledger still lists `lemma_16_10` as open: the false predicate is preserved for audit, and the numbered result awaits a faithful contextual repair. The intake re-derived every step of Theorem II.233.1 against the Lean definitions at the pin and checked the cover envelope exhaustively on $\mathbb{F}_7$ and $\mathbb{F}_{11}$; source 50’s two scripts were rerun on copies at the write, with identical output.

F.56 Source 51. The conclusion and its quantifiers

Source 51, delivered Section 1.

Added in the ninth write, 7 October 2026. Delivered title block. “An obstruction and a repair for Section 13 square geometry / Report279 / Research report prepared for Vladimir / 7 October 2026”; *the PDF author field reads “Research report prepared for Vladimir”; no statement on AI assistance is made.* The title page is dated 7 October 2026; the archive arrived at 20:20 on 6 October 2026 (Pacific time). Archive Report279_Section13_Square_Geometry_Obstruction_and_Repair.zip (arrival 5d4718772), manuscript article.tex (948 lines, 22 pages), six pins from aad4789fc (definitions and budgets) to 4f8cd1198 (its SOURCES.md and provenance/source_manifest.json, shipped as 51-square-geometry-SOURCES.md and data/51-square-geometry-source_manifest.json; the eighteen Lean snapshots it ships are not staged). It is the fifth report of the series of sources 26, 32, 41 and 42 (Report275–278). Printed below: the abstract and Section 1; Sections 2–10 stand in the Part II chapter on Section 13 after source 53 (Section II.195), Sections 11–12 and Appendix A are in Part V. Its macros `\code`, `\len` ($\ell(P)$, one argument), `\st` (d_P) and `\pin` are renamed `\codefiftyone`, `\lenfiftyone`, `\stfiftyone` and `\pinfiftyone`. It numbers its equations consecutively through the manuscript; here they are numbered within sections. Its appendix opened a small-type group and a `samepage` environment before the appendix heading; the group now opens after the heading, and `samepage` is dropped.

Abstract of source 51 (as delivered). For arbitrarily large prime moduli we construct a genuine Section 13 context, complete Stage 13.4 and 13.5 data, and explicit Stage 13.6–13.9 data satisfying all six numerical budgets in the pinned bilinear-extraction theorem and the retained upper length bound $|R| \leq |Q|$. Nevertheless, the chosen product $S \times U$ contains no proper common-step square of side length two. Thus these hypotheses and retained predicates do not imply the source’s exact square-grid predicate for arbitrary witnesses. The construction uses a horizontal strip, the zero function, a uniformly controlled Fourier spectrum, and two progressions with disjoint nonzero difference sets. Its density $\alpha = \lfloor N/8 \rfloor / N$ varies with N ; every asymptotic estimate is uniform on $[1/16, 1/8]$.

This refutes a universal retained-data implication, not existential extraction, particular proof-term witnesses, Corollary 13.10, or Szemerédi’s theorem. We also prove a constructive counterpart for suitably selected witnesses. Preserving the short parent from the actual budget construction, a balanced aligned partition and short-carrier rectification connect the existing all-scale coefficient theorem to the already available quarter-density cover. The direct selection attains the original effective density and width constants while bypassing the unchanged Stage 13.9 interface. The geometric proofs are self-contained; the analytic coefficient theorem is explicitly imported from pinned source. The final compared source already constructs a complete square with one large-prime threshold uniform above any fixed positive lower density bound; we distinguish this from its older fixed-exact-density API. No new Lean verification is claimed. The companion uses bounded exact diagnostics and curated source hashes, never astronomical numerical witnesses.

The question concerns information retained between successive extraction steps. A source theorem can construct suitable witnesses without its recorded output predicates characterizing every witness that would work in a later proof. Confusing these two assertions is consequential here: the Stage 13.9 step relation is modular, whereas a square-cover argument needs a quantitative relation between ordinary integer indices.

We use the source definitions at commit aad4789f of the public repository *VladimirReshetnikov/ProveIt*. The stage definitions, explicit budget theorem, and exact-grid definition are unchanged at the later commit 330102ae. That later commit adds a conditional common-step covering theorem, discussed in Section II.210. Later pinned sources add aligned parent selection, complete fixed-density square extraction, and finally a threshold uniform over varying actual densities; Sections II.210.3–II.210.6 record those developments. The complete file map and byte-verification scope appear in Appendix V.197.

A modular progression is an indexed object

$$P(a, d; L) = \{a + jd : 0 \leq j < L\} \subseteq \mathbb{Z}/N\mathbb{Z}.$$

It is *proper* when its displayed parameterization is injective, equivalently when its carrier has cardinality L . Modular wrapping is permitted. Nonzero step is a separate requirement in the source, including for singleton progressions. We write $\ell(P) = L$ and $d_P = d$. For proper progressions we freely identify their lengths with their carrier cardinalities.

Theorem F.56.1 (Full-budget failure of the universal grid implication). *For every sufficiently large prime N , there exist:*

- (i) *a `Section13Context` with $\alpha = \lfloor N/8 \rfloor / N \in [1/16, 1/8]$, $\eta = 2^{-44}$, and $\phi = 0$;*
- (ii) *data D, E satisfying the complete predicates `IsStage134Data` at the cutoff $\theta = \lambda_{10}(\alpha^{32}/16)$ and `IsStage135Data`;*
- (iii) *a positive integer m satisfying all six explicit budgets of the theorem `section13_bilinear_extraction_of_budgets`;*
- (iv) *data F, G, H, J satisfying all four predicates `IsStage136Data`, `IsStage137Data`, `IsStage138Data`, and `IsStage139Data`, as well as $\ell(R) \leq \ell(Q)$;*

for which $y = 0$, the final retained set is exactly $S \times U$, and

$$(S - S) \cap (U - U) = \{0\}. \quad (\text{F.56.1})$$

Consequently $S \times U$ contains no product of two proper progressions of the same nonzero step and length at least two. In particular, `Corollary1310Grid` is false for these selected data.

Section II.211 proves the positive counterpart: suitable witness selection under the same budgets yields a bilinear square with the intended effective constants. The negative theorem gives a counterexample to the universal assertion

$$\text{full antecedent and all retained predicates} \implies \text{grid for those same retained witnesses}.$$

It leaves the following existential statements untouched:

- The budget theorem asserts the existence of F, G, H, J satisfying its output predicates. Our chosen data themselves satisfy those predicates.
- The newer context-to-extraction theorem constructs a whole chain at each fixed positive density at most $1/6$, once the prime modulus is sufficiently large. We neither rerun nor replace that construction.
- Section II.211 arranges additional geometry by suitable reselection. The counterexample does not identify the witnesses chosen by any particular Lean proof term.
- The catalogue conclusion of Corollary 13.10 does not require its output square to lie inside $S \times (U + y)$, or its output set to lie inside $J.D$. Since $\phi = 0$ everywhere here, that conclusion is readily satisfied by a full step-one square.

These distinctions are part of the claim, rather than qualifications added after it.

[write] **7 October 2026; sources 47–53.** *Identifications and credit within this report.* Source 51 (Report279) continues the series of sources 26, 32, 41 and 42 (Report275–278) and cites none of them; sources 52 and 53 (Report280 and Report281) are of the same day and cite neither this source nor each

other. Source 53 repeats this source’s rectification and balanced-parent lemmas and its endpoint-contained cover (Lemmas II.201.1, II.201.2, II.201.3) and sharpens its repair (note after Corollary II.211.4). Its six pins (aad4789fc, 330102ae0, 2548e7fff, 38e74b4f0, 77f7d8c1f, 4f8cd1198) are ancestors of HEAD, and the 24 rows of `data/51-square-geometry-source_manifest.json` match the snapshots it ships (not staged here); the companion was rerun on a copy at the write, normally and with `-O` (identical diagnostics; the source-snapshot step needs its Linux builder and was replaced by the provenance check).

Since the pins (checked at HEAD 06b30ca52, ledger 113 exact companions and 7 open statements). The repository has repaired the Section 13 interfaces so that geometry constructed by the preceding lemmas is carried forward (8da5b7f07, 7 October 2026): `corollary_13_10` (`Sections12_13.lean`, line 686) is an exact companion, `corollary_13_10_holds` (`Proofs13ConstructedChain.lean`, line 31), consuming the step-span and construction certificates exported by Lemmas 13.7 and 13.9; the older arbitrary-witness interface is preserved, unasserted, as `corollary_13_10_without_construction_geometry` (`Sections12_13.lean`, line 669), and `Corollary1310Grid` (`Proofs13BilinearRestriction.lean`, line 33) with the grid-qualified `corollary_13_10_of_stage_chain_and_grid` (line 270) remain. Theorem F.56.1 concerns that older retained interface; `FORMALIZATION_STATUS.txt`: “Source 51’s grid obstruction concerns the older retained data, not the newly proved construction certificates.” Neither the obstruction nor this source’s repair is formalized as stated; the repository’s construction follows the short-parent route with its own exponents. The file and line ranges cited in this source refer to its pins; of the files it cites, `Sections12_13.lean`, `Proofs13BilinearExtraction.lean` and `Proofs13CommonStepSelection.lean` have changed since.

F.57 Source 52. Main conclusions and source scope

Source 52, delivered Section 1.

Added in the ninth write, 7 October 2026. Delivered title block. “Composite modulus obstructions to Stage 13.7 / Report280 / Research report prepared for Vladimir / 7 October 2026”; *the PDF author field reads “Research report prepared for Vladimir”; no statement on AI assistance is made.* The title page is dated 7 October 2026; the archive arrived at 20:20 on 6 October 2026 (Pacific time). Archive `Report280_Composite_Modulus_Obstructions_to_Stage137.zip` (arrival 5d4718772), manuscript `article.tex` (835 lines, 20 pages), pin 5c9a442d2 (its `SOURCES.md` and `provenance/source_manifest.json`, shipped as `52-composite-modulus-SOURCES.md` and `data/52-composite-modulus-source_manifest.json`; the five Lean snapshots it ships are not staged). It is the sixth report of the series of sources 26, 32, 41, 42 and 51. Printed below: the abstract and Section 1; Sections 2–9 stand in the Part II chapter on Section 13 after source 51 (Section II.195), Sections 10–11 and Appendices A–B are in Part V. Its macro `\code` (the same definition as source 51’s) is renamed `\codefiftyone`. It numbers its equations consecutively through the manuscript; here they are numbered within sections.

Abstract of source 52 (as delivered). The preserved propositions `lemma_13_7_without_domain` and `lemma_13_7_without_prime_assumption` are false at the pinned source interface. We give exact finite counterexamples: a short progression of even residues in $\mathbb{Z}/2^k\mathbb{Z}$ supports an order-eight Freiman map given by division of its integer coordinates by two, while every ambient affine function agrees with it at at most one point. The Stage 13.7 length and mass requirements are therefore incompatible. Every field of the context and every Stage 13.6 conjunct is verified directly, with the source’s unnormalized arrangement count.

A second explicit construction also satisfies the genuine Stage 13.4 and 13.5 predicates. A further variant makes the input progressions have unit step; ambient affine agreement is then bounded by two points per row. Thus adding the initial-stage predicates alone, or requiring a unit common input step alone, does not repair the statements. The explicit unit-step construction satisfying both initial-stage predicates has $\ell R > \ell Q$. A separate nonunit-step family removes that limitation: it has genuine initial stages at the exact Section 10 cutoff, $R \subseteq Q$, and all six numerical budgets of the source extraction theorem, while still admitting no Stage 13.7 output. The prime hypothesis of that theorem is false for these examples. The elementary obstruction extends to every fixed

prime-power family, and a two-point model in $\mathbb{Z}/18\mathbb{Z}$ explains why bounded-order Freiman identities need not imply ambient affine extendability. These are mathematical refutations of unproved encoded propositions, not counterexamples to the prime-corrected Lean theorem, a new Lean verification, a priority claim, or a claim about the global Szemerédi argument. The companion performs bounded exact checks without constructing the astronomical moduli.

The key distinction is between affine dependence on a progression parameter and affine dependence on the ambient residue. Over a field these descriptions are interchangeable when the progression step is nonzero. Over a composite modulus that conversion can fail, even when all short additive relations are preserved. In the present interface it is the ambient meaning that is required.

All source claims in this report are pinned to the public repository *VladimirReshetnikov/ProveIt*, commit

5c9a442d263632fdcd4e9690154c12c2dcc70b53.

The definitions were freshly read at this pin. The complete byte contents of `Definitions.lean` and `Sections12_13.lean` agree with the earlier construction pin 77f7d8c1; the hashes and precise source ranges are recorded in Appendix V.200. We do not track unrelated changes after this fixed source comparison.

Theorem F.57.1 (An explicit counterexample with minimal input requirements). *Let*

$$k = 939 \cdot 2^{1892} + 5, \quad N = 2^k, \quad m = N/16.$$

There are a genuine `Section13Context` S_0 of density $\alpha = 1/16$ and records D, E, F such that:

- (i) *`IsStage135Data` and `IsStage136Data` hold;*
- (ii) *$D.P = E.Q = F.R = (0, 2, m)$, so $\ell R = \ell Q$;*
- (iii) *no record G satisfies `IsStage137DataWithoutDomain`.*

In particular both preserved all-moduli Stage 13.7 propositions are false. No Stage 13.4 predicate is asserted for this first construction.

Here “minimal input requirements” describes the logical premises being used. It is not a claim that N is the smallest counterexample modulus.

Theorem F.57.2 (Genuine initial stages and unit input steps). *There is a second, explicitly specified dyadic modulus with a context of density $1/16$, an allowed spectral threshold $\theta = \alpha^{32}/4$, and records D, E, F satisfying all three predicates*

`IsStage134Data`, `IsStage135Data`, `IsStage136Data`,

for which no weak Stage 13.7 output exists. This can be done with $D.P$, $E.Q$, and $F.R$ all of step 1. In this genuine-initial-stage construction,

$$\ell(D.P) = \ell(E.Q) = 5 < \ell(F.R).$$

Separately, there is an explicit unit-step counterexample with $D.P = E.Q = F.R$ and genuine Stage 13.5 and 13.6 predicates, without a Stage 13.4 certificate.

Theorem II.217.1 further gives a nonunit-step family with genuine initial stages, $R \subseteq Q$, the exact Section 10 cutoff and all six numerical budgets. Theorem F.57.2 deliberately separates its two sets of extra properties. It does not refute a statement requiring *simultaneously* the genuine Stage 13.4 predicate, a unit common step, and $\ell R \leq \ell Q$.

What is refuted. The two target declarations are `def ... : Prop`. They universally quantify over a nonzero natural modulus, a context and records, then assume only `IsStage136Data`. Their conclusions assert the existence of G satisfying respectively the weak and full Stage 13.7 predicates. A record called `Stage134Data` does not contain a certificate of `IsStage134Data`; the latter is a separate predicate. The same distinction applies at Stage 13.5. These quantifiers are material to Theorem F.57.1, although Theorem F.57.2 removes that particular concern by checking both extra predicates directly.

What remains separate. The corrected proposition `lemma_13_7` quantifies over N with Fact N.Prime. The pinned file `Proofs13CompleteRowExtraction.lean` contains the theorem `lemma_13_7_holds : lemma_13_7`. Our moduli are composite and cannot instantiate that hypothesis. We have read its statement and source, but have not rerun its Lean dependency build or kernel audit. Nothing here challenges that corrected result, later prime-modulus square extraction, or a correlation theorem that takes ambient polynomial coefficients as part of its input.

[write] **7 October 2026; sources 47–53.** *Identifications within this report.* Source 52 (Report280) continues the series of sources 26, 32, 41, 42 and 51 and cites none of them; source 51 and source 53 are of the same day. Its pin `5c9a442d2` is an ancestor of HEAD; the five rows of `data/52-composite-modulus-source_manifest.json` match the snapshots it ships (not staged here). The companion was rerun on a copy at the write, normally and with `-0` (identical diagnostics; the source-snapshot step needs its Linux builder and was replaced by the provenance check). The intake checked the order-eight Freiman property of $\phi(2j) = j$ and the agreement bound one by exhaustion for $N = 2^k$, $k = 4, \dots, 8$, and the $\mathbb{Z}/18\mathbb{Z}$ model.

Formal status (checked at HEAD `06b30ca52`, ledger 113 exact companions and 7 open statements). The two refuted propositions are still preserved and unproved: `lemma_13_7_without_domain` and `lemma_13_7_without_prime_assumption` (Sections12_13.lean, lines 526 and 534); `FORMALIZATION_STATUS.txt` calls them “preserved and unproved”, and no negation is formalized. They are not ledger statements, and the ledger is unchanged by this source. The prime-corrected `lemma_13_7` (line 558) is an exact companion; since the pin its statement also exports integer step-span geometry, its proof is `lemma_13_7_holds` (`Proofs13AllScalesStepSpan.lean`, line 118), and the theorem at lines 46–53 of `Proofs13CompleteRowExtraction.lean` cited below is now `lemma_13_7_without_step_span_holds` (`8da5b7f07`). Nothing here bears on either, since every modulus here is composite. The line ranges of the source map below refer to the pin; `Sections12_13.lean`, `Proofs13BilinearExtraction.lean` and `Proofs13CompleteRowExtraction.lean` have changed since.

F.58 Source 53. The explicit theorem

Source 53, delivered Section 1.

Added in the ninth write, 7 October 2026. Delivered title block. “Explicit uniform density square extraction with a sharper side length / Report281 / Research report prepared for Vladimir / 7 October 2026”; the PDF author field reads “Research report prepared for Vladimir”; no statement on AI assistance is made. The title page is dated 7 October 2026; the archive arrived at 20:20 on 6 October 2026 (Pacific time). Archive `Report281_Explicit_Uniform_Density_Square_Extraction.zip` (arrival `5d4718772`), manuscript `article.tex` (756 lines, 16 pages), pin `bf2fe146d` (its `SOURCES.md` and `provenance/source_manifest.json`, shipped as `53-explicit-threshold-SOURCES.md` and `data/53-explicit-threshold-source_manifest.json`; the twenty-three Lean snapshots it ships are not staged). It is the seventh report of the series of sources 26, 32, 41, 42, 51 and 52, and the spine of the Section 13 chapter. Printed below: the abstract and Section 1; Sections 2–8 open the Part II chapter on Section 13 (Section II.195), Sections 9–10 and Appendix A are in Part V. Its macros `\code`, `\st` (the operator step) and `\pin` are renamed `\codefiftythree`, `\stfiftythree` and `\pinfiftythree`. It numbers its equations consecutively through the manuscript; here they are numbered within sections.

Abstract of source 53 (as delivered). We give a single explicit threshold for bilinear square extraction from the prime-modulus Section 13 context, uniformly over all actual densities above a prescribed positive lower bound. If $0 < a_0 \leq a \leq 1$, the threshold depends only on a_0 , while both the side exponent and retained density use the actual a . The seven sufficient scale and rounding conditions collapse exactly to

$$N_0(a_0) = \left\lceil \left(\frac{8}{c(a_0)} \right)^{1/\gamma(a_0)} \right\rceil.$$

Here c and γ are explicit positive functions given below. At this same threshold, suitable choices in the all-density construction give a proper common-step square of side at least $\frac{1}{2}N^{4\gamma(a)} - 1 \geq N^{3\gamma(a)}$, retaining density at least $2^{-137}a^{704}$ and agreement with the original phase on the retained set.

The geometric improvement preserves the actual short parent in the Stage 13.6 witness, rectifies two subprogressions, uses balanced aligned parents, and retains the full coefficient-localization exponent. The scalar threshold comparison and geometric lemmas are proved here; the analytic extraction theorems are explicitly imported from pinned source. The basic density-uniform square theorem is already upstream. This report supplies a mathematical quantitative refinement, without claiming a new Lean implementation, kernel replay, priority result, arbitrary-witness catalogue repair, or global Szemerédi bound. Bounded exact diagnostics accompany the proof and never evaluate the astronomical thresholds.

For $0 < a \leq 1$, define

$$K(a) = 2^{114}a^{-320}, \quad Q(a) = 2^{1048576}a^{-2097152}, \quad (\text{F.58.1})$$

$$c(a) = \frac{1}{2}2^{-228K(a)}a^{576K(a)}, \quad \beta(a) = 2^{-13Q(a)}, \quad (\text{F.58.2})$$

$$f(a) = 2^{-100}a^{448}, \quad g(a) = 2^{-284}a^{1408}, \quad (\text{F.58.3})$$

$$\gamma(a) = \frac{\beta(a)f(a)g(a)}{4} = 2^{-386}a^{1856}2^{-13Q(a)}. \quad (\text{F.58.4})$$

All real powers in this report have positive bases. Natural lengths, floors and ceilings are regarded as real numbers when used in such powers.

Theorem F.58.1 (Explicit uniform extraction with actual density). *Fix $0 < a_0 \leq 1$ and put*

$$N_0(a_0) = \left\lceil \left(\frac{8}{c(a_0)} \right)^{1/\gamma(a_0)} \right\rceil. \quad (\text{F.58.5})$$

For every prime $N \geq N_0(a_0)$ and every Section 13 context (A, ϕ, a) with $a_0 \leq a \leq 1$, there are proper progressions V, W in $\mathbb{F}_N$ with the same nonzero step and the same length k , and a set B , such that

$$k \geq D(a)N^{4\gamma(a)} - 1 \geq \frac{1}{2}N^{4\gamma(a)} - 1 \geq N^{3\gamma(a)}, \quad (\text{F.58.6})$$

$$B \subseteq A \cap (V \times W), \quad |B| \geq 2^{-137}a^{704}k^2, \quad (\text{F.58.7})$$

where

$$D(a) = 2^{-g(a)}c(a)^{f(a)g(a)} > \frac{1}{2}. \quad (\text{F.58.8})$$

There are coefficients $c_{00}, c_{10}, c_{01}, c_{11} \in \mathbb{F}_N$ such that

$$\phi(x, z) = c_{00} + c_{10}x + c_{01}z + c_{11}xz \quad ((x, z) \in B). \quad (\text{F.58.9})$$

The context assumptions are stated in Section II.196. In particular, the theorem does not concern an arbitrary dense set with an arbitrary phase.

[write] **7 October 2026; sources 54–59.** *It makes source 57’s threshold explicit* (Section II.287): since γ increases with the density, this theorem gives the conclusion of the uniform square extraction used by source 57 for every prime above $\lceil (8/c(\delta))^{1/\gamma(\delta)} \rceil$, and so a closed threshold for source 57’s Theorem F.62.1 (note after its Remark II.290.4). The two sources do not know each other.

The coefficient-bearing bound in (F.58.6) is the more informative estimate. The factor three in the clean exponent is a convenient consequence at the stated threshold; it is not claimed to be optimal.

Corollary F.58.2 (A simpler larger threshold). *Theorem F.58.1 also holds with*

$$N_{\text{safe}}(a_0) = \left\lceil 2^{2^{14 \cdot 2^{1048576} a_0^{-2097152}}} \right\rceil \quad (\text{F.58.10})$$

in place of $N_0(a_0)$.

For an expanded formula containing only the input a_0 , let $h_0 = \log_2(1/a_0)$. Then (F.58.5) is exactly

$$\left\lceil 2^{[4 + 2^{114} a_0^{-320} (228 + 576 h_0)] 2^{386 + 13 \cdot 2^{1048576} a_0^{-2097152}} a_0^{-1856}} \right\rceil. \quad (\text{F.58.11})$$

The expression has one final rounding operation. It is intended as an exact symbolic threshold, not as a numerically feasible size.

Relationship with existing source. The pinned repository already proves a theorem with one threshold for every context above a fixed positive density bound, using $\gamma(a_0)$ and $2^{-137} a_0^{704}$ in its exported conclusion. Its proof retains actual-density mass before the last weakening. The refinement here makes the threshold explicit, keeps the actual density in the exponent, and strengthens the length by using suitable short-parent witnesses. It does not claim that the basic uniformity theorem is missing.

All repository comparisons are frozen at

`bf2fe146d323cf2828ca0de63c9a965ca0638eaa.`

The source map in Appendix V.204 identifies the imported results and records the scope of the comparison. The proof below is an ordinary mathematical derivation from these inputs, not a newly compiled formal proof.

[write] **7 October 2026; sources 47–53.** *Identifications and credit within this report.* Source 53 (Report281) is the seventh report of the series of sources 26, 32, 41, 42, 51 and 52, cites none of them, and is the spine of this report’s Section 13 chapter (Section II.195). *Its geometric part is source 51’s*, of the same day: Lemmas II.201.1, II.201.2 and II.201.3 are source 51’s Lemmas II.211.1, II.211.2 and Proposition II.210.1 with proofs in nearly the same words, and Proposition II.201.4 is source 51’s Theorem II.211.3 without its final weakening $\lceil L/2 \rceil^e \geq L^{e/2}$, which is why the exponent here is $4\gamma(a)$ against source 51’s $2\gamma(\alpha)$ (note after source 51’s Corollary II.211.4); the quantifier observation of Section II.196 is source 51’s (Section II.210.5). New here: the explicit threshold and its exact collapse, the actual-density exponent and mass, and the coefficient $D(a) > 1/2$. The pin `bf2fe146d` is an ancestor of HEAD; the 23 rows of `data/53-explicit-threshold-source_manifest.json` match the snapshots it ships (not staged here); the companion was rerun on a copy at the write, normally and with `-0` (identical diagnostics; the snapshot step needs the Linux builder).

Formal status (checked at HEAD `06b30ca52`, ledger 113 exact companions and 7 open statements). *Since the pin the repository has made its own threshold explicit:* `section13_complete_square_extraction_explicit` (Proofs13ExplicitGeometricThreshold.lean, line 24; commit 93d436414, 7 October

2026, 06:14) proves, for every prime $N \geq \text{section13GeometricThreshold } \delta$ (line 12, a finite maximum of rounded recurrence, integer and square-scale thresholds) and every context with $\alpha \geq \delta$, a common-step square of side at least $N^{\gamma(\delta)}$ and mass $2^{-137}\delta^{704}$. So “the threshold is existential” is no longer true at HEAD; the theorem above remains stronger in three respects that are not formalized: the single closed threshold $N_0(a_0)$ with its exact collapse (Proposition II.199.1), the actual density in the exponent and in the mass, and the side $D(a)N^{4\gamma(a)} - 1 \geq N^{3\gamma(a)}$. The repository’s later fixed-exact-density theorems reach larger exponents, but with the quantifier form criticized in Section II.196 (note in source 51’s Section II.210.5). `theorem_13_12` is an exact companion (`theorem_13_12_holds`, `Proofs13EndpointPrintedTheorem.lean`, line 73); the uniform theorem of the pin is `section13_complete_square_extraction_uniform_density` (`Proofs13UniformSquareExtraction.lean`, line 15).

F.59 Source 54. The result and its place in the proof

Source 54, delivered Section 1.

Added in the tenth write, 7 October 2026. Delivered title block. “A Single Quadratic Phase Across All Translates / Spectral extraction, offset rigidity, and sharp stability in Gowers’s Section 8 / Research manuscript prepared for the ProveIt project / 6 October 2026”; *the PDF author field reads the same; no statement on AI assistance is made.* Archive `Gowers_Single_Phase_Research.zip` (arrival `aa1d86f6f`), manuscript `article.tex` (1,387 lines, 20 pages), pin `6a6d7961c` (its `SOURCE_AUDIT.md`, shipped as `54-single-phase-SOURCE_AUDIT.md`, records four blob identifiers and the README blob `b2e25fa`). Printed below: the abstract and Section 1; Sections 2–9 form a chapter of Part II on Proposition 8.1 (Section II.267), Sections 10–12 are in Part V. Its macro `\ee` (the letter e, not this report’s one-argument `e(\cdot)`) is set as this report’s `\e`, which prints the same; its ‘question’ environment (“Question”, on the shared counter) is set in this report’s ‘Question’ style.

Abstract of source 54 (as delivered). We strengthen the quadratic-bias extraction in Proposition 8.1 of Gowers’s 2001 proof of Szemerédi’s theorem. One global quadratic polynomial suffices simultaneously on every translate of the input set; no arithmetic-progression hypothesis is needed for the extraction itself. For the balanced function of a set of density δ , the coefficient $1/\sqrt{2}$ can be replaced by $[\delta(1 - \delta)\max(\delta, 1 - \delta)]^{-1}$, which is at least $27/4$. The underlying statement is an exact mixed, weighted Fourier factorization with a spectral-overlap parameter and optimal coefficient one. Quadratic refinements extend it to every finite abelian group, including even-order groups with genuinely nonclassical quadratic phases.

We also identify the exact global energy ceilings for an affine selector’s intercept: after removing its quadratic part, an offset of order 1, 2, 3, or at least 4 has normalized ceiling 1, $1/2$, $1/3$, or $1/4$, respectively. We classify all equality cases and prove dimension-free stability. Orders at most four have the optimal L^2 recovery exponent $1/2$; every order at least five has the optimal exponent $1/4$. A four-frequency construction proves that the latter loss cannot be removed. A graph formulation extends the ceilings to simultaneous intercepts via the classical Motzkin–Straus theorem. All assertions are proved in the article; finite checks and an integration plan accompany it. These are local structural improvements, not a claim of an improved end-to-end Szemerédi bound.

F.59.1 The Section 8 interface

Let N be odd, let $A \subseteq \mathbb{Z}/N\mathbb{Z}$ have density $0 < \delta < 1$, and write $f = \mathbf{1}_A - \delta$. Put

$$\sigma = \mathbb{E}_x |f(x)|^2 = \delta(1 - \delta), \quad B = \|f\|_\infty = \max(\delta, 1 - \delta). \quad (\text{F.59.1})$$

Here and throughout, group averages are normalized. For comparison with Gowers, define the *unnormalized* derivative Fourier coefficient

$$D_h(r) = \sum_{x \in \mathbb{Z}/N\mathbb{Z}} f(x) \overline{f(x-h)} e(-rx/N), \quad e(t) = \exp(2\pi it). \quad (\text{F.59.2})$$

The conjugation does not change this expression for a balanced function, which is real-valued.

Proposition 8.1 of [1] starts with an arithmetic progression P of cardinality T and an affine selector $h \mapsto \lambda h + \mu$ satisfying

$$\sum_{h \in P} |D_h(\lambda h + \mu)|^2 \geq \beta N^2 T. \quad (\text{F.59.3})$$

It produces quadratic polynomials ψ_s , allowed to depend on the translate, for which

$$\sum_s \left| \sum_{z \in P+s} f(z) e(-\psi_s(z)/N) \right| \geq \frac{\beta N T}{\sqrt{2}}. \quad (\text{F.59.4})$$

The corrected repository statement explicitly records oddness of N ; its proof module is `Proofs08QuadraticBias.lean` [223].

Theorem F.59.1 (One common quadratic polynomial). *Let N be odd and let $\emptyset \neq P \subseteq \mathbb{Z}/N\mathbb{Z}$ be arbitrary. Assume $\beta > 0$ and (F.59.3). There is one $r \in \mathbb{Z}/N\mathbb{Z}$ such that the polynomial*

$$\psi(z) = 2^{-1} \lambda z^2 + r z \quad \text{in } \mathbb{Z}/N\mathbb{Z} \quad (\text{F.59.5})$$

satisfies

$$\sum_s |f(s)| \left| \sum_{z \in P+s} f(z) e(-\psi(z)/N) \right| \geq \frac{\beta N T}{\sigma}, \quad (\text{F.59.6})$$

$$\sum_s \left| \sum_{z \in P+s} f(z) e(-\psi(z)/N) \right| \geq \frac{\beta N T}{B\sigma} \geq \frac{27}{4} \beta N T. \quad (\text{F.59.7})$$

More precisely, σ in both denominators can be replaced by the spectral overlap $\mathcal{D}_\mu(g)$ defined in Section II.268, where $g(z) = f(z) e(-2^{-1} \lambda z^2/N)$.

[write] 7 October 2026; sources 54–59. *Against Proposition 8.1 and the formal statement* (checked at the write). For $0 < t = \min(\delta, 1 - \delta) \leq 1/2$, $B\sigma = t(1 - t)^2$ has derivative $(1 - t)(1 - 3t)$, so its maximum is $4/27$ at $t = 1/3$ and $1/(B\sigma) \geq 27/4 > 1/\sqrt{2}$. Taking $\psi_s = \psi$ for every s therefore gives Gowers’s conclusion, with a polynomial of degree two in z (the inverse 2^{-1} exists because N is odd), for every nonempty P , so in particular for the progression of the formal statement. `proposition_8_1` (`Sections08_09.lean`, line 29) assumes N odd and a progression P (`ModAP`) and asks for translate-dependent polynomials ψ_s of degree at most two with the constant $1/\sqrt{2}$; for densities 0 and 1 the balanced function vanishes and its hypothesis forces $\beta \leq 0$. The theorem thus implies the catalogue statement on paper; the catalogue statement is already exact (`proposition_8_1_holds`, `Proofs08QuadraticBias.lean`, line 1069), and nothing of this source is formalized. The constant $27/4$ is the uniform consequence of the amplitude bound, not an optimum over balanced indicators (Remark after the proof; item II.57).

The same ψ occurs in *every* summand. The theorem does not assert a lower bound on every individual translate, nor does it assert that the same polynomial works for every choice of P . Its improvements over the cited interface are the common phase, the amplitude dependence, the explicit

constant, and the removal of the progression hypothesis. The last point concerns extraction only: later recurrence and partition arguments may still require progression geometry.

The theorem follows from a short Fourier factorization, but that factorization also retains information normally lost by pointwise Cauchy–Schwarz. We develop its mixed form, its equality deficit, and a constructive selector before studying its extremal consequences.

F.59.2 Offset rigidity and its stability

For a finite abelian group G , let $\widehat{G}$ be its character group, written additively as homomorphisms $G \rightarrow \mathbb{T} = \mathbb{R}/\mathbb{Z}$. For $g : G \rightarrow \mathbb{C}$ define

$$A_\eta(g)(h) = \mathbb{E}_x g(x+h) \overline{g(x)} e(-\eta(x)), \quad \mathcal{E}_\eta(g) = \mathbb{E}_h |A_\eta(g)(h)|^2. \quad (\text{F.59.8})$$

Let $\sigma = \|g\|_2^2 > 0$. The exact identity proved below is

$$\frac{\mathcal{E}_\eta(g)}{\sigma^2} = \sum_{\xi \in \widehat{G}} p_\xi p_{\xi-\eta}, \quad p_\xi = \frac{|\widehat{g}(\xi)|^2}{\sigma}. \quad (\text{F.59.9})$$

Thus the relevant extremal problem is a quadratic program on disjoint cycles. Its value is a specialization of classical graph optimization [224]; the explicit stability statement is developed here.

Theorem F.59.2 (Exact offset ceilings and recovery exponents). *Put $m = \text{ord}(\eta)$ and*

$$c_m = \begin{cases} 1, & m = 1, \\ 1/2, & m = 2, \\ 1/3, & m = 3, \\ 1/4, & m \geq 4. \end{cases} \quad (\text{F.59.10})$$

Then $\mathcal{E}_\eta(g) \leq c_m \sigma^2$, and every constant is best possible. All extremizers are described by the power-spectrum table in Theorem II.271.3.

If $0 \leq \varepsilon \leq 1/16$ and $\mathcal{E}_\eta(g) \geq (c_m - \varepsilon) \sigma^2$, there is an exact extremizer g_ with $\|g_*\|_2^2 = \sigma$ such that*

$$\|g - g_*\|_2 \leq \begin{cases} K_m \sqrt{\sigma} \varepsilon^{1/2}, & m \leq 4, \\ 3\sqrt{\sigma} \varepsilon^{1/4}, & m \geq 5, \end{cases} \quad (K_1, K_2, K_3, K_4) = (\sqrt{2}, \sqrt{5}, \sqrt{10}, 3). \quad (\text{F.59.11})$$

The exponents are optimal, uniformly over the ambient group whenever the indicated nontrivial families exist. The constants K_m and 3 are not claimed optimal.

For $m \geq 5$, an extremizer has at most three Fourier frequencies in a single η -orbit, with half the power at the central frequency and half on its two neighbors. Near an intersection of two such extremal families, a small fourth frequency costs only the square of its power. This is the source of the fourth-root recovery exponent, not a defect of a repeated Cauchy–Schwarz argument.

F.59.3 Scope and attribution

The ingredients used here—finite Fourier inversion, quadratic refinements, elementary graph quadratic programs, and conversion from power-spectrum distance to L^2 distance—are standard. No

novelty claim is made for them or for the Motzkin–Straus extremal value. The proposed contribution is their precise common-phase formulation at the Section 8 interface, the resulting constants, the full offset classification, and the explicit sharp-exponent recovery theorem and localization consequences. The inspected repository index and source audit were checked for overlap; the entire research corpus was not exhaustively audited. Priority in the wider literature is therefore not certified.

A local improvement must not be confused with the current global quantitative problem. For example, Leng–Sah–Sawhney prove $r_k(N) \ll N \exp(-(\log \log N)^{c_k})$ for $k \geq 5$ [2]. The September 2026 preprint of Milićević gives a high-characteristic quasipolynomial inverse theorem via Freiman multihomomorphisms [32]. Neither result is a proof dependency here, and neither is superseded by the present local extraction.

[write] **7 October 2026; sources 54–59.** *Identifications and credit within this report.* [223] is this repository at 6a6d7961c, an ancestor of HEAD; the blobs recorded in 54-single-phase-SOURCE_AUDIT.md (the transcription, Sections08_09.lean 0fb6e0a, source 36’s source audit) are unchanged at HEAD fe7165a3c, and the “research index” it read in selected ranges is this directory’s README (blob b2e25fa). Its “corrected repository statement” is proposition_8_1 (Sections08_09.lean, line 29), with the exact companion proposition_8_1_holds (Proofs08QuadraticBias.lean, line 1069) (note after Theorem F.59.1). *Not cited:* source 31, whose phase-space function contains the offset energy of this source (note after Lemma II.271.1); source 31 was placed but not yet written in at this source’s pin, and the README did not yet describe it, so this is an overlap, not a lapse. This is the report’s first chapter on Gowers’s Section 8 (Section II.267). Source 54’s suite was rerun on a copy at the write.

F.60 Source 55. The endpoint problem and its place in Gowers’s proof

Source 55, delivered Section 1.

Added in the tenth write, 7 October 2026. Delivered title page. “ProveIt research manuscript / Sharp Second-Order Stability at the Polynomial-Phase Endpoint of Gowers Norms / A parity-sensitive refinement of the Szemerédi framework / Prepared with ChatGPT for Vladimir Reshetnikov / October 2026”; *the PDF author field reads “ChatGPT; research manuscript prepared for Vladimir Reshetnikov”.* Archive gowers_endpoint_stability.zip (arrival aa1d86f6f), manuscript article.tex (1,498 lines, 21 pages), pin 128f514af (its SOURCE_AUDIT.md, shipped as 55-endpoint-stability-SOURCE_AUDIT.md, also names the earlier commit 6a6d7961c). Printed below: the abstract, the status statement of its title page and Section 1; Sections 2–11 form, with source 57’s Section 5 after them, a chapter of Part II on the Proposition 6.1 endpoint (Section II.236), Sections 12–13 and Appendices A–B are in Part V. It uses cleveref (options nameinlink, noabbrev); its references are written out as cleveref prints them (lowercase names). Its macro $\mathcal{U}$ ($\mathcal{U}_d$, not this report’s $\mathcal{U}$) is renamed $\mathcal{U}_{fiftyfive}$; its ‘question’ environment (“Question”, on the shared counter) is set in this report’s ‘Question’ style.

Abstract of source 55 (as delivered). Let $d \geq 2$, let G be a nontrivial finite abelian group, and let $|f| \leq 1$. Write $\varepsilon = 1 - \|f\|_{U^d(G)}^{2^d}$ and let $D_d(f)$ be the squared normalized L^2 distance from f to the phase polynomials of degree at most $d - 1$, including nonclassical phases. We prove

$$D_d(f) \leq 2^{1-d} \varepsilon + b_d(G) \varepsilon^2 + O_d(\varepsilon^3), \quad b_d(G) = \begin{cases} (3^d - 1)/(2 \cdot 4^d), & |G| \text{ odd}, \\ (1 - 4^{-d})/3, & |G| \text{ even}. \end{cases}$$

The error and the sufficiently small range are uniform in G . Both coefficients are sharp on every fixed nontrivial group of the indicated parity. We also prove a torsion-profile refinement and classify the tangent perturbations of second-order extremizers: a conjugate pair of characters in odd order, and a single order-two character in even order.

The local estimate has a fully explicit radius and cubic error. Its central ingredient is a fifth-order cancellation for bounded complex functions, obtained from triple independence of cube vertices and the pointwise constraint $|f| \leq 1$. A four-vertex lattice census identifies the two-torsion contribution. The passage from near-maximal norm to the explicit local neighborhood uses the published qualitative theorem of Eisner and Tao. No new global bound in Szemerédi’s theorem is asserted.

Status. AI-assisted research draft with complete mathematical proofs and reproducible finite checks. It has not been refereed or verified in a proof assistant. Publication priority for the sharpened statements is not established. The distinction between imported results, proved refinements, and proposed research questions is maintained throughout.

Gowers’s proof of Szemerédi’s theorem organizes the analysis of arithmetic patterns around iterated multiplicative differences, uniformity, and the extraction and integration of structured phase information [1]. Its small-uniformity counting estimates and its inverse steps are complementary: the former control pseudorandom errors, while the latter locate algebraic structure. This article concerns the extreme structured end of that framework, where a normalized Gowers norm is close to its maximum.

This is not the general inverse problem at a fixed small positive norm. Here the objective is sharper: determine the best possible distance to an exact polynomial phase as the cubical defect tends to zero. The resulting estimate can be applied after an inverse or integration argument has already produced a sufficiently accurate phase candidate. It does not, by itself, improve the density range in the full Szemerédi iteration.

F.60.1 Normalization and the question answered

All expectations below are normalized counting averages. Put $q = 2^d$, and for a function $f : G \rightarrow \mathbb{C}$ define

$$\mathcal{U}_d(f) = \|f\|_{U^d(G)}^q = \mathbb{E}_{x, h_1, \dots, h_d \in G} \prod_{\omega \in \{0,1\}^d} \mathcal{C}^{|\omega|} f(x + \omega \cdot h), \quad (\text{F.60.1})$$

where $\mathcal{C}$ is complex conjugation and $|\omega|$ is the sum of the coordinates. The quantity in (F.60.1) is real and nonnegative. For example, this follows by induction from the multiplicative-difference recursion and $\|f\|_{U^1}^2 = |\mathbb{E}f|^2$. If $|f| \leq 1$, then $0 \leq \mathcal{U}_d(f) \leq 1$.

Definition F.60.1. The set $\mathcal{P}_{d-1}(G)$ consists of the maps $P : G \rightarrow S^1$ for which

$$\prod_{\omega \in \{0,1\}^d} \mathcal{C}^{|\omega|} P(x + \omega \cdot h) = 1 \quad \text{for every } x, h_1, \dots, h_d \in G.$$

Define

$$D_d(f) = \min_{P \in \mathcal{P}_{d-1}(G)} \mathbb{E}|f - P|^2, \quad \varepsilon(f) = 1 - \mathcal{U}_d(f). \quad (\text{F.60.2})$$

Thus “degree” refers to multiplicative differences, not to a chosen coordinate formula. In small characteristic the definition includes nonclassical phases. Replacing this class by classical polynomials would change the problem and would invalidate the general statement. The minimum exists: $\mathcal{P}_{d-1}(G)$ is a closed subset of the finite product $(S^1)^G$, hence compact. Multiplication by any of its elements preserves $\mathcal{U}_d$.

Our principal question is the asymptotic behavior of

$$\mathfrak{S}_{d,G}(e) = \sup\{D_d(f) : |f| \leq 1, \varepsilon(f) \leq e\}. \quad (\text{F.60.3})$$

It will be determined through order e^2 , with a uniform $O_d(e^3)$ error.

F.60.2 Published input and the repository comparison

Eisner and Tao’s compact-group near-extremizer theorem states that $\|f\|_{U^d} \rightarrow 1$, under $|f| \leq 1$, forces L^1 distance to a polynomial phase to tend to zero, uniformly in the group [7, Theorem 1.1]. Their Remark 1.6 notes effective polynomial rates without optimizing or displaying them. We use only the qualitative entry statement. Everything after entry into an explicitly specified neighborhood is proved directly here.

The inspected ProveIt collection already contains several local cube expansions, torsion-sensitive counting estimates, moment inequalities, and rigidity statements for additive energy [227]. In particular, the fourth-order arithmetic of cube configurations is not presented as a wholly new idea. The additions developed here are the uniform sharp distance expansion, the bounded-complex fifth-order cancellation that makes its error cubic, the order-two spectral profile, and the associated second-order rigidity theorem. The contemporary work of Kovač and Rogers on L^p variants and near-extremizers [71] was also consulted; this manuscript addresses the specific normalized L^∞ endpoint distance problem (F.60.3).

The comparison is with an inspected repository snapshot, not a claim to have refereed every earlier manuscript. Searches and the primary sources consulted are recorded in `SOURCE_AUDIT.md`. No claim that the results are the first published formulations of these constants is needed for their proofs.

[write] **7 October 2026; sources 54–59.** *Identifications and credit within this report.* [227] is this repository at `128f514af`, an ancestor of HEAD (its `55-endpoint-stability-SOURCE_AUDIT.md` also names `6a6d7961c`); the “several local cube expansions, torsion-sensitive counting estimates, moment inequalities, and rigidity statements” it inspected are sources of this report (among them sources 08, 31, 38 and 45). *Not cited:* sources 31 and 38, whose sharp endpoint coefficients at $d = 3$ are this source’s $d = 3$ coefficients in another class of functions (note after Theorem II.237.1); both were placed but not yet written in at this source’s pin, and the README did not yet describe their endpoint results, so this is an overlap, not a lapse. [7] is the report’s entry for the Eisner–Tao paper (source 55 cites its journal version; sources 02 and 57 the preprint). Source 55’s suite was rerun on a copy at the write.

F.61 Source 56. The question and the main result

Source 56, delivered Section 1.

Added in the tenth write, 7 October 2026. Delivered title block. “Nearly Sharp Relative Affine Repair on Fourier-Uniform Domains / Weighted Freiman tests, error-independent uniformity, and explicit quadratic-level-set consequences / Research draft prepared for the ProveIt project / 6 October 2026”; *the PDF author field reads “Research draft for ProveIt”*. Archive `gowers_relative_affine_repair.zip` (arrival `aa1d86f6f`), manuscript `article.tex` with thirteen section files under `sections/` (1,593 lines in all, 25 pages; the main file `\inputs` the sections, which are printed here in that order), pin `b6a8ee966` (itself the arrival commit of sources 47–50), with `128f514af` recorded as a later search head “not used as pin” (its `source_manifest.json`, shipped as `data/56-relative-affine-repair-source_manifest.json`). Printed below: the abstract, the status statement after it and Section 1; Sections 2–9 stand in the Part II chapter on affine rigidity after source 59 (Section II.248), Sections 10–12 and Appendix A are in Part V. Its macro `\Q` (calligraphic $\mathcal{Q}$, not this report’s $\mathbb{Q}$) is set as source 49’s `\Qcal`, the same definition; its ‘question’ environment (“Question”, on the shared counter) is set in this report’s ‘Question’ style.

Abstract of source 56 (as delivered). Let G be a finite abelian group, let H be an arbitrary abelian group, and let $w : G \rightarrow [0, 1]$ have mean $\alpha > 0$. We prove an explicit relative stability theorem for the additive-quadruple test. If $\|w - \alpha\|_{U^2(G)}^4 \leq \alpha^5/1024$ and a map on the support

of w violates an $\varepsilon \leq 1/1024$ fraction of the weighted additive quadruples, then it agrees with a unique affine map except on weighted relative mass at most $13\varepsilon/50$. The uniformity requirement is independent of ε . A two-sided error profile improves the coefficient to $1/4 + o(1)$ as both the normalized uniformity parameter and ε tend to zero; $1/4$ is optimal. The proof retains autocorrelation weights through modal difference selection, uses an exact triangle second-moment identity, applies classical homomorphism self-correction, and then counts quadruples with exactly one error. For a quadratic level set cut out by r equations over $\mathbb{F}_p$, p odd, the same conclusion follows when every nonzero quadratic combination has matrix rank R with $p^{R-4r} \geq 2^{14}$. This is an affine-repair conclusion on a quadratic domain, not a quadratic-repair theorem. Exact rational verification code, finite tests, a source comparison, and a formalization plan accompany the proofs. Linear repair itself is classical; the contribution proposed here is the explicit weighted relative theorem and its quantitative refinements.

Status. Complete mathematical proofs are given below. This is an AI-generated research manuscript, not externally refereed or proof-assistant verified. Computational tests are diagnostics, not proofs of the universal statements. Publication priority beyond the sources inspected is not asserted.

The approximate-homomorphism stages of Gowers’s proof convert information about many respected additive configurations into genuine algebraic structure [1]. A useful strengthening in the high-agreement regime should not merely find a small structured restriction. It should repair almost the whole original domain. On an arbitrary sparse domain this is false: the additive relations may fail to determine any global affine extension. The question studied here is whether a quantitatively uniform domain makes that repair possible with an error proportional to the observed violation rate, without tightening the uniformity assumption every time the violation rate decreases.

Our answer is affirmative under an explicit U^2 condition. All averages in this article are probability averages. In particular, the fourth power of a U^2 norm contains no hidden factor of $|G|$.

F.61.1 Weighted tests and distance

Let G be a finite abelian group of order N , and let H be an abelian group, written additively. No finiteness assumption on H is needed. Fix $w : G \rightarrow [0, 1]$ with

$$\alpha = \mathbb{E}_x w(x) > 0, \quad u = \|w - \alpha\|_{U^2(G)}.$$

A map ϕ need only be specified on $\text{supp } w$; elsewhere it can be extended arbitrarily. Define the weighted configuration mass and defect by

$$\mathcal{Q}(w) = \mathbb{E}_{a,b,c \in G} w(a)w(b)w(c)w(a+b-c), \quad (\text{F.61.1})$$

$$\mathcal{D}_w(\phi) = \mathbb{E}_{a,b,c \in G} w(a)w(b)w(c)w(a+b-c) \mathbf{1}_{\phi(a)+\phi(b) \neq \phi(c)+\phi(a+b-c)}. \quad (\text{F.61.2})$$

Thus the actual test error is

$$\varepsilon = \frac{\mathcal{D}_w(\phi)}{\mathcal{Q}(w)}. \quad (\text{F.61.3})$$

The denominator is positive because $\mathcal{Q}(w) \geq \alpha^4$. Every tuple is ordered, and repeated vertices are included. An affine map $\Phi : G \rightarrow H$ means $\Phi(x) = L(x) + c$, with $L \in \text{Hom}(G, H)$ and $c \in H$. Its relative error on the weighted domain is

$$t_w(\phi, \Phi) = \frac{\mathbb{E}_x w(x) \mathbf{1}_{\phi(x) \neq \Phi(x)}}{\alpha}. \quad (\text{F.61.4})$$

For $w = \mathbf{1}_A$, these definitions are exactly the fraction of additive quadruples in A that are violated and the fraction of points of A that must be changed.

Theorem F.61.1 (Explicit weighted relative affine repair). *Suppose that*

$$u^4 \leq \frac{\alpha^5}{1024}, \quad 0 \leq \varepsilon \leq \frac{1}{1024}. \quad (\text{F.61.5})$$

There is a unique affine map $\Phi : G \rightarrow H$ satisfying

$$t_w(\phi, \Phi) \leq \frac{13}{50}\varepsilon. \quad (\text{F.61.6})$$

It is obtained by three finite modal selections: first of differences on weighted pairs, then of differences of the selected map on G , and finally of the constant residual on the original domain.

Uniqueness in this statement means uniqueness *within the displayed error radius*. It does not assert that G admits only one affine map. An externally supplied upper bound for the test error can replace the actual ε in the theorem. The proof first obtains the weaker but simple estimate $t \leq \varepsilon/3$, and then improves it to (F.61.6).

Corollary F.61.2 (Subset formulation). *Let $\emptyset \neq A \subseteq G$ have density α . If*

$$\|\mathbf{1}_A - \alpha\|_{U^2(G)} \leq 2^{-5/2}\alpha^{5/4}$$

and $\phi : A \rightarrow H$ respects at least a $1 - \varepsilon$ fraction of the additive quadruples in A , where $0 \leq \varepsilon \leq 2^{-10}$, then ϕ agrees with a unique affine map on at least $(1 - 13\varepsilon/50)|A|$ points of A .

The condition is a hypothesis on the domain, not on the labels. In particular, it does not involve ε . It applies in every characteristic, on cyclic groups of composite order, and on products of finite abelian groups.

F.61.2 The nearly sharp profile

The useful scale-free quantities are

$$\lambda = \frac{u^2}{\alpha^{5/2}}, \quad \kappa = \frac{u^4}{\alpha^4} = \lambda^2\alpha, \quad s = \frac{u^2}{\alpha^2} = \lambda\sqrt{\alpha}. \quad (\text{F.61.7})$$

The main hypothesis is $\lambda \leq 1/32$. The following statement holds without that hypothesis and without an upper bound on ε .

Theorem F.61.3 (Two-sided local error profile). *For every affine map Φ , put $t = t_w(\phi, \Phi)$. Then*

$$\varepsilon(1 + \lambda^2\alpha) \geq 4(1 - \lambda\sqrt{1 - \alpha})t - 12t^2 - 12\lambda t^{3/2}, \quad (\text{F.61.8})$$

$$\varepsilon(1 + \lambda^2\alpha) \leq 4(1 + \lambda\sqrt{1 - \alpha})t. \quad (\text{F.61.9})$$

For the repaired map in Theorem F.61.1, one has in addition

$$t \leq \frac{(1 + \lambda^2\alpha)\varepsilon}{4(1 - \lambda\sqrt{1 - \alpha}) - 4\varepsilon - 4\sqrt{3}\lambda\sqrt{\varepsilon}}. \quad (\text{F.61.10})$$

The denominator in (F.61.10) is positive under (F.61.5).

Consequently, along any family with $\varepsilon > 0$, $\varepsilon \rightarrow 0$, and $\lambda \rightarrow 0$, the repaired distance satisfies $t/\varepsilon \rightarrow 1/4$. A single changed label on a large group proves that $1/4$ cannot be replaced by a smaller limiting coefficient. The density exponent $5/4$, the threshold $1/1024$, and the finite-error coefficient $13/50$ are *not* asserted to be optimal.

F.61.3 Relation to existing results and limits of the claim

Linear repair is not new. Kazhdan–Ziegler’s polynomial splining theorem already gives $C\varepsilon$ repair on sufficiently uniform subsets of finite-field vector spaces, with a condition of the form $u < (\alpha\varepsilon)^B$ [228, Theorem 1.7]. Their exact extension theorem is also relevant. Peluse’s recent elementary relative affine lemma gives error $16\varepsilon^{1/4}$ under $u < (\alpha\varepsilon)^{32}$ and $0 < \varepsilon < 2^{-15}$ [229, Lemma 7.4]. In that overlapping regime, Corollary F.61.2 has both weaker uniformity requirements and a smaller error bound. This comparison is with that local lemma, not with Peluse’s global inverse theorem.

The proposed contribution is the explicit error-independent condition $u^4 \leq \alpha^5/1024$, the weighted formulation on arbitrary finite abelian domains, the quantitative profile, and the consequences proved here. The classical self-correction step is attributed to the Blum–Luby–Rubinfeld method [49]. Earlier qualitative repair results can sometimes be bootstrapped to an error-independent hypothesis; we do not claim that this phenomenon itself has no precedent.

Within the inspected ProveIt material, the existing full-domain affine energy profiles and local Bohr/Freiman extension results address different hypotheses [230]. The present theorem is not a replacement for those statements on all their inputs. It supplies a new high-agreement branch when the domain has the required Fourier uniformity. No improved global bound for Szemerédi’s theorem, no improved general Balog–Szemerédi exponent, and no new global U^4 inverse bound is deduced in this article.

[write] **7 October 2026; sources 54–59.** *Identifications and credit within this report.* [230] is this repository at `b6a8ee966`, an ancestor of HEAD; `Sections06_07.lean`, `Section10.lean` and `Proofs10Shift.lean` have their pin blobs at HEAD `fe7165a3c`, and the README it read is blob `b2e25fa`. The “existing full-domain affine energy profiles and local Bohr/Freiman extension results” are those of sources 11, 12 and 14 (Sections II.22 and II.33); their high-agreement statements concern the whole group, this theorem a uniform sparse domain. Source 33’s question on sparse repair of arrangements (item II.25) is the closest earlier question; this theorem answers its affine, quadruple-test analogue, not the arrangement version. The comparisons with Kazhdan–Ziegler and Peluse are this source’s and were not re-read at the write. Source 56’s suite (`verify.py`, `repair.py`) was rerun on a copy at the write.

F.62 Source 57. Principal results and their scope

Source 57, delivered Section 1.

Added in the tenth write, 7 October 2026. Delivered title block. “Quantitative Bilinear Extraction and Bounded Additive Defects / Refinements of Gowers’s argument and the ProveIt Ramsey research programme / Prepared for the ProveIt Ramsey project / 7 October 2026”; *the PDF author field reads “Research draft prepared with ChatGPT for the ProveIt project”*. The title page is dated 7 October 2026; the archive arrived at 21:12 on 6 October 2026 (Pacific time). Archive `gowers_research_proveit.zip` (arrival `aa1d86f6f`), manuscript `gowers_bilinear_extraction_and_additive_defects.tex` (2,340 lines, 31 pages), pin `128f514af` (its `provenance.json`, shipped as `data/57-bilinear-extraction-provenance.json`, records six Lean blobs, the transcription, `FORMALIZATION_STATUS.txt` and the archive blobs of sources 40 and 45 at `8ed496b14`). Printed below: the abstract, the repository-baseline statement after it and Section 1; Sections 2–4 form a Part II chapter on Theorem 13.12 (Section II.287), Section 5 stands in the Part II chapter on the Proposition 6.1 endpoint after source 55 (Section II.236), Section 6 forms a Part II chapter on energy levels (Section II.291), Sections 7–8 are in Part V. It uses `cleveref` (options `capitalise`, `noabbrev`); its references are written out with the name of the statement referred to, capitalized. *Its delivered PDF names every statement on the shared counter “Theorem”* (for example “Theorem 2.2” for Lemma 2.2 and “Theorem 4.4” for Remark 4.4), because its `cleveref` identifies those environments only by their counter; the names printed here are those of the statements. Its ‘Input theorem’ environment (on the shared counter) is added to this report; its ‘question’ environment (“Research question”, numbered 1–12 on its own counter) is set in this report’s ‘Research question’ style, numbered within its section, which keeps the indices.

Abstract of source 57 (as delivered). We develop quantitative and structural refinements of the framework in Gowers’s 2001 proof of Szemerédi’s theorem, using a pinned snapshot of the ProveIt Ramsey development as the starting point. First, we combine an earlier Fejér arrangement-selection theorem with the repository’s uniform common-square extraction. Keeping exact intermediate exponents reduces the prepared density loss to $(\alpha/2)^{2^{996}201429}$. An elementary near-maximal U^4 argument completes the range near $\alpha = 1$. Together these give a written proof of the printed Theorem 13.12 length bound $p^{2^{-\alpha-2^{70}}}$, with the retained-density exponent improved from 2^{76} to 2^{41} . We also prove explicit near-maximal frequency-extraction estimates at every uniformity order, a finite abelian group extension, and an explicit cubic-phase approximation bound in characteristic at least five.

Second, if a positive real n -set has Schur-triple deficit s and $n \geq 5s - 1$, we prove that it is a positive dilation of $\{j + \lambda_j : 1 \leq j \leq n\}$, where the nonnegative nondecreasing integers λ_j sum to s . Thus there are exactly $p(s)$ eventual dilation classes. A sharper treatment of deficit two classifies all third-level additive-energy extremizers for $m \geq 7$, with the optimal threshold and its six-point exception. These answer two explicit questions in the preceding repository research.

All added results have written proofs. We identify the inherited inputs precisely, supply exact-arithmetic checks, and propose a research programme. The manuscript does not claim a new final Szemerédi bound, a new Lean proof, or established publication priority.

Repository baseline. `128f514afce0bf7eb48133fb30dd3f8e034ea3a2` (7 October 2026 UTC). This is an AI-assisted research manuscript prepared for mathematical review and repository integration.

Gowers’s argument [1] converts a large uniformity norm into structured Fourier information, then uses that structure in a density increment. Two places in this framework invite a precise improvement. The first is quantitative: a loss in the density of an intermediate frequency graph appears later inside a very large exponent. The second is structural: extremal additive energy records how an ordered set can fail to complete additive configurations. We treat both questions, with the proof dependencies kept separate.

F.62.1 A strengthened Section 13 extraction theorem

For the statement below, $\widehat{f}(r) = p^{-1} \sum_x f(x) \exp(-2\pi i r x / p)$ and $\Delta_h f(x) = f(x + h) \overline{f(x)}$.

Theorem F.62.1 (Printed length, stronger retained density). *For every $0 < \alpha < 1$ there is an integer $N_0(\alpha)$ such that, for every prime $p \geq N_0(\alpha)$ and every $f : \mathbb{F}_p \rightarrow \mathbb{C}$ with $|f| \leq 1$ and $\|f\|_{U^4}^{16} > \alpha$, the following objects exist: proper arithmetic progressions $P, Q \subseteq \mathbb{F}_p$, with the same nonzero common difference and the same integer length m ; a set $B \subseteq P \times Q$; and coefficients $c_{00}, c_{10}, c_{01}, c_{11} \in \mathbb{F}_p$. They satisfy*

$$m \geq p^{2^{-\alpha-2^{70}}}, \quad (\text{F.62.1})$$

$$|B| \geq (\alpha/2)^{2^{41}} m^2, \quad (\text{F.62.2})$$

$$\left| \widehat{\Delta_h \Delta_k f}(c_{00} + c_{10}h + c_{01}k + c_{11}hk) \right| \geq \alpha/2 \quad ((h, k) \in B). \quad (\text{F.62.3})$$

In Gowers’s unnormalized convention, the last bound is $\alpha p/2$. The length in (F.62.1) is the printed Theorem 13.12 target. The stronger mass in (F.62.2) replaces its exponent 2^{76} by 2^{41} . The statement uses the sufficiently-large-prime quantifier order already present in the pinned Lean catalogue. The case $\alpha = 1$ in that catalogue is vacuous because its nonuniformity hypothesis is strict.

The inspected development already proves a complete geometric square extraction, but its older input density led to an exponent explicitly proved smaller than the printed target [232]. We resolve that *written-proof gap relative to the pinned snapshot*. The classical theorem itself is Gowers’s; our quantitative improvement and proof route are the additions.

Put $a = \alpha/2$ and $d = 2\,996\,201\,429$. The comparison is:

Quantity	Previous inspected route	This article
Density supplied to the square extraction	$a^{2^{66}}$	a^d , and hence at least $a^{2^{32}}$
An explicit length valid for every $0 < \alpha < 1$	$p^{2^{-(2/\alpha)^{2^{88}}}}$	$p^{2^{-(2/\alpha)^{2^{53}}}}$
Retained density on the square	$a^{2^{76}}$	$a^{2^{41}}$
Near-maximal range	Not supplied by that propagation	Whole ambient square; good-pair density at least $1 - 332(1 - \alpha)$ when $1 - \alpha \leq 2^{-10}$

The proof is modular. Gowers’s original Fourier and Freiman inputs produce a separately Freiman frequency graph. The earlier Fejér selection theorem in repository source 40 [233], reproved in the required specialization below, improves its purification loss. The existing density-uniform geometric theorem then supplies actual proper progressions. A separate elementary near-maximal argument covers the interval near one where the factors α and $\alpha/2$ cannot be interchanged inside exponentials. The detailed size threshold is recorded in Remark II.290.4.

F.62.2 Explicit rigidity near maximal uniformity

The endpoint argument yields more than the preceding application requires. If $Q_4(f) := \|f\|_{U^4}^{16} \geq 1 - \varepsilon$ and $\varepsilon \leq 2^{-10}$, Theorem II.247.5 supplies a scalar $c \in \mathbb{F}_p$ and a set of at least $(1 - 332\varepsilon)p^2$ pairs with

$$\left| \widehat{\Delta_h \Delta_k f}(chk) \right| \geq \frac{3}{4} - 4\varepsilon.$$

The argument combines a dimension-independent phase-normalization estimate, elementary additive correction, and the exact symmetry of mixed derivatives. It requires no inverse theorem as an input.

For s derivative parameters, Theorem II.247.7 replaces 332 by $216s + 8$ and the frequency chk by $ch_1 \cdots h_s$. The finite abelian group version Theorem II.247.9 supplies an s -additive frequency map, with an explicit recurrence for its error constant. In characteristic at least five the cubic frequency can be integrated: Corollary II.247.10 gives, for $0 \leq \varepsilon < 1/17604$, a polynomial P of degree at most three with

$$|\mathbb{E}_x f(x) e_p(-P(x))| \geq \sqrt{1 - 6194\varepsilon} - \varepsilon.$$

The corresponding squared L^2 error is at most 12390ε . More generally, Corollary II.247.11 gives polynomial-phase approximation of degree at most $k - 1$ when $p \geq k$, with explicit constant $K_k = 2 \prod_{r=1}^{k-2} (108r + 1)$ multiplying the uniformity deficit. The numerical constants are not asserted to be optimal, but the linear order of the squared error cannot be improved uniformly. Eisner–Tao’s general near-extremizer theorem [7] is the appropriate earlier context; the present results make a particular elementary route and its constants explicit.

F.62.3 A partition theorem and a complete energy equality case

For a positive real set $C = \{c_1 < \cdots < c_n\}$, define

$$h(C) = \binom{n}{2} - \#\{(x, y, z) \in C^3 : x + y = z\}.$$

Both orders of a pair of distinct summands are counted, and equal summands are permitted. The maximum Schur count and its equality case are classical [231]. We prove the following exact classification beyond a linear threshold:

$$h(C) = s, \quad n \geq 5s - 1 \quad \Longleftrightarrow \quad C = u\{j + \lambda_j : 1 \leq j \leq n\},$$

where $u > 0$, the integers λ_j are nonnegative and nondecreasing, and $\sum_j \lambda_j = s$. There are therefore exactly $p(s)$ positive dilation classes. The perturbed positions lie among the last s entries. This is Theorem II.292.3; the threshold is sufficient, with its optimal value left open for general s .

For $s = 2$, Theorem II.292.4 improves the threshold to the optimal $n = 6$. This sharpened endpoint classification identifies every equality case at the third energy level. Let

$$E(A) = \#\{(a, b, c, d) \in A^4 : a + b = c + d\}, \quad M_m = (2m^3 + m)/3.$$

For $m \geq 7$, the m -element real or torsion-free sets satisfying $E(A) = M_m - 8(m - 3)$ are exactly the affine images of

$$J_m = \{0, 1, \dots, m\} \setminus \{2\}, \quad K_m = \{0, 2, 3, \dots, m - 1, m + 1\}.$$

At $m = 6$ there is precisely one further affine class, $\{0, 1, 2, 4, 5, 6\}$. Thus seven is the optimal threshold for the two-class statement. The earlier numerical energy levels and the second-level classification are credited to repository source 45 [234]; the endpoint recurrence has earlier literature precedents [189]. The new classification answers the first two research questions posed in that source.

F.62.4 What the results establish

The Fourier statements concern prime cyclic groups, with a separate finite abelian group near-maximal extension. The energy statements concern ordered real sets and their torsion-free transfers. Moving the latter to finite cyclic groups requires a suitable Freiman isomorphism or rectification theorem; arbitrary integer representatives do not suffice.

The final density-increment iteration is not carried out here. Consequently these local refinements establish no new upper bound for $r_k(N)$, the largest size of a subset of $[N]$ with no k -term progression. For perspective, Leng, Sah, and Sawhney [2] prove $r_k(N) \ll N \exp(-(\log \log N)^{c_k})$ for $k \geq 5$ by a different modern inverse-theorem route. This article makes no comparison claim against that global bound.

Every added theorem below has a written proof. The exact finite computations are supplementary checks, not proofs of unbounded families. The substantial imported geometric theorem is identified by its declaration and commit. No new Lean module, repository status change, or fresh compilation of the full formal development is claimed. The literature check establishes relevant precedents and repository-relative additions; publication priority remains unestablished.

[write] **7 October 2026; sources 54–59.** *Identifications within this report.* [232] is this repository at 128f514af, an ancestor of HEAD; of the six Lean blobs recorded in data/57-bilinear-extraction-provenance.json five are unchanged at HEAD fe7165a3c, while Sections12_13.lean has changed (its definitions of theorem_13_12, secondDifferenceFourier and BilinearOn are unchanged; theorem_13_12 has moved from line 626 to line 718), and so have the transcription and FORMALIZATION_STATUS.txt (from d846374). [233] and [234], read from their archives at 8ed496b14 before their placement, are this report's source 40 (Section F.45) and source 45 (Section F.50); this source's numbering "source 40", "source 45" agrees with the report's. Its Fejér selection is source 40's Theorem II.169.2 with the feature analysis of Theorem II.166.2, specialized and reproved (Section II.289); its energy recurrence and first two levels are source 45's (Section II.183). *Since the pin*, theorem_13_12 has been proved formally along this source's route (note after Theorem II.290.3). Both verifiers were rerun on copies at the write.

F.63 Source 58. Main results and their place in the existing work

Source 58, delivered Section 1.

Added in the tenth write, 7 October 2026. Delivered title block. "Polynomial restriction on every cyclic group / A torsion-sensitive refinement of Gowers's Lemma 9.3 and an optimal exponent for positive Fourier filtering / Research contribution prepared for the ProveIt project / 7 October 2026"; *the PDF author field reads "Research contribution prepared for ProveIt"; its Appendix B says that the report "was prepared with OpenAI ChatGPT"*. The title page is dated 7 October 2026; the archive arrived at 21:12 on 6 October 2026 (Pacific time). Archive gowers_torsion_uniform_restriction.zip (arrival aa1d86f6f), manuscript gowers_torsion_uniform_restriction.tex (1,770 lines, 26 pages), pin 5d4718772 (the arrival commit of sources 51–53; its SOURCE_AUDIT.txt and source_provenance.json, shipped as 58-torsion-restriction-SOURCE_AUDIT.txt and data/58-torsion-restriction-source_provenance.json). Printed below: the abstract, its keywords and Section 1; Sections 2–11 form a chapter of Part II on Lemma 9.3 (Section II.276), Section 12 to Appendix B are in Part V. Its macros \GG (sans-serif G, not this report's $\Gamma_{\delta,s}$), \commit (its pin) and \repo (in small capitals) are renamed \GGfiftyeight, \commitfiftyeight and \repsc; its unnumbered 'Theorem A' and 'Theorem B' environments are added to this report; its 'question' environment ("Question", numbered within sections) is set in this report's 'Question' style; its figure is shipped as figures/58-torsion-restriction-kernel_exponents.pdf.

Abstract of source 58 (as delivered). Gowers's restriction lemma extracts a set on which almost all additive sixteen-tuples are respected by a given map, while retaining a controlled

number of configurations. We prove a version for every cyclic modulus and arbitrary finite abelian target in which the retention exponent is 54, in place of 2^{19} , with the explicit sufficient hypothesis $|B| \geq (4/(\alpha\eta))^{54}$. The proof amplifies a fixed positive Fourier filter and accounts explicitly for both small torsion and repeated entries. For a general balanced relation of length m , the central estimate has nonprincipal error $\tau_G(2D)(R_m^k - P^k)|B|^{m-2}$ and collision error $\binom{m}{2}R_2^k|B|^{m-2}$. It yields polynomial restriction with exponent $O(m \log m)$ on every cyclic group.

We also determine the optimal leading exponent within the normalized nonnegative Fourier-filter class. If a nonzero target discrepancy may have fixed order $a \geq 2$, then $\inf_K [-\log P_m(K)/\log(P_m(K)/Q_{m,a}(K))] = m \log m / (2 \log a) + O_a(m \log \log m)$. The lower bound uses an elementary coefficient rearrangement inequality; powered cosine filters attain the leading term. All statements count ordered tuples with repetitions. The article includes full proofs, exact integer certificates, a source comparison, integration interfaces, and a research programme. These are local restriction results; no stronger global Szemerédi bound or completed Lean formalization is asserted.

Keywords. Quantitative additive combinatorics; approximate homomorphisms; Fourier restriction; finite abelian groups; small torsion; Szemerédi’s theorem.

The quantitative inverse step in Gowers’s proof of Szemerédi’s theorem requires more than a large subset of a graph. It requires a subset on which almost every additive configuration has an additive image, while enough configurations survive to feed the subsequent argument. Lemma 9.3 of [1] supplies such a restriction for sixteen entries. Its retained configuration proportion is

$$\left(\frac{\alpha\eta}{4}\right)^{2^{19}}. \quad (\text{F.63.1})$$

The exponent is a local cost in a much longer proof. Reducing it does not, on its own, determine the final density-increment bound.

We use the following notation throughout. For $m = 2r \geq 4$, a subset B of a finite abelian group G , and a map $\phi : B \rightarrow H$ to another finite abelian group, let $\mathsf{T}_m(B)$ be the number of ordered tuples in B^m with equal sums of the first and last r entries. Let $\mathsf{G}_m(B, \phi)$ count those tuples for which the corresponding sums of the images are also equal. Every count includes repeated entries. Write $b = |B|$.

Theorem A (Sixteen-variable restriction on every cyclic group). *Let $N \geq 1$, let $B \subseteq \mathbb{Z}/N\mathbb{Z}$, and let $\phi : B \rightarrow H$, where H is any finite abelian group. Suppose $0 < \alpha, \eta \leq 1$ and*

$$b \geq \left(\frac{4}{\alpha\eta}\right)^{54}, \quad \mathsf{G}_{16}(B, \phi) \geq \alpha b^{15}. \quad (\text{F.63.2})$$

Then there is $B' \subseteq B$ such that

$$\mathsf{G}_{16}(B', \phi) \geq (1 - \eta)\mathsf{T}_{16}(B'), \quad (\text{F.63.3})$$

$$\mathsf{T}_{16}(B') \geq \left(\frac{\alpha\eta}{4}\right)^{54} b^{15}. \quad (\text{F.63.4})$$

Consequently $|B'| \geq (\alpha\eta/4)^{54/15}|B|$.

[write] **7 October 2026; sources 54–59.** *Checked at the write* (exact integers). With $A = \sum_{|n| \leq 25} \binom{50}{25+n}^{16}$ and B its even- n part: $A \geq 2^{750}$, $250A \geq 481B$, $481^{53} \geq 2^{50}250^{53}$, $481^{50} \geq 2^{47}250^{50}$, $\binom{100}{50} \leq 2^{97}$,

$4(50 + \binom{16}{2}) = 680$, $680 \cdot 2^{97}/2^{108} = 85/256 < 1$ and $2^{108}/2^{105} = 8$ (Appendix V.220). *Against the formal statement.* `lemma_9_3` (Sections08_09.lean, line 77) quantifies over all $N \geq N_0$, $|B| = \beta N$ and maps $\mathbb{Z}/N\mathbb{Z} \rightarrow \mathbb{Z}/N\mathbb{Z}$, with input $\alpha\beta^{15}N^{15} \leq \text{phiAdditiveTupleCount } 8$ (that is, $\mathbf{G}_{16}(B, \phi) \geq \alpha\beta^{15}$; for nonempty B it forces $\alpha \leq 1$) and asks for $(\alpha\eta/4)^{2^{19}}\beta^{15}N^{15} \leq \text{additiveTupleCount } 8$ and $\text{GammaHomOfOrder } 8 \ B' \ \phi \ (1 - \eta)$; with $N_0 = \lceil \beta^{-1}(4/(\alpha\eta))^{54} \rceil$ this theorem gives both, since $0 < \alpha\eta/4 < 1$ and $54 < 2^{19}$. So it implies `lemma_9_3` on paper with an explicit threshold; the catalogue statement is already exact (`lemma_9_3_holds`, `Proofs09Restriction.lean`, line 376), and this theorem is not formalized.

[write] 7 October 2026; sources 60–69. *Exponent 53 by source 68* (Section II.310): the same statement with 53 in place of 54, through a degree-twelve rational filter (its Theorem II.312.2); for cyclic targets source 62's powers $\alpha^{31}\eta^{30}$ are better when $\alpha^{22}\eta^{23}(1 - \eta) < 2^{-139}$ (note in that chapter).

The map in this statement need not be a homomorphism on B , and N may have any small prime factors. The sufficient threshold is explicit but large. The integer 54 is a convenient certified exponent, not a claimed best constant. Section II.283 gives a more precise version whose retained powers are approximately $\alpha^{53.2064}\eta^{52.2064}$, at a threshold with power approximately 49.3412 in $(\alpha\eta)^{-1}$.

The proof separates two phenomena. Small torsion in the *target* controls how many principal Fourier frequencies an incorrect relation can pass. Small torsion in the *domain* controls how many solutions an additional, nonprincipal integer relation can have. The former is handled by a spectral gap; the latter by the kernel sizes of multiplication maps. Repeated tuples require a third estimate, because selecting a point once is different from selecting each of its occurrences independently.

The general expectation estimate, proved in Theorem II.282.1, keeps all three costs explicit. A filter of degree D has nonnegative symmetric decreasing coefficients c_n ; we set

$$P = \sum_n c_n^m, \quad Q = \sum_{2|n} c_n^m, \quad R_s = \int_{\mathbb{T}} K(t)^s dt.$$

After k independent layers, good tuples have expected survival at least P^k . The bad-tuple expectation is at most

$$Q^k \mathbf{U}_m(B, \phi) + \left\{ \tau_G(2D)(R_m^k - P^k) + \binom{m}{2} R_2^k \right\} b^{m-2}, \quad (\text{F.63.5})$$

where $\mathbf{U}_m = \mathbf{T}_m - \mathbf{G}_m$ and $\tau_G(L) = \max_{1 \leq d \leq L} |\ker(d : G \rightarrow G)|$. For a cyclic group, $\tau_G(2D) \leq 2D$, independently of its modulus.

There is also a sharp asymptotic result about the method itself. For a normalized nonnegative trigonometric polynomial with nonnegative Fourier coefficients, define

$$P_m(K) = \sum_n c_n^m, \quad Q_{m,a}(K) = \sum_{a|n} c_n^m, \quad \mathcal{A}_{m,a}(K) = \frac{-\log P_m(K)}{\log(P_m(K)/Q_{m,a}(K))}.$$

The ratio is assigned $+\infty$ if $P_m = Q_{m,a}$.

Theorem B (Optimal leading exponent in the positive-filter class). *For each fixed integer $a \geq 2$, as the integer m tends to infinity,*

$$\inf_K \mathcal{A}_{m,a}(K) = \frac{m \log m}{2 \log a} + O_a(m \log \log m). \quad (\text{F.63.6})$$

The infimum is over finite real even trigonometric polynomials satisfying $0 \leq K \leq 1$, $K(0) = 1$, and $c_n \geq 0$; monotonicity of the coefficients is not required for the lower bound. Powered cosine filters attain the displayed leading term.

[write] **7 October 2026; sources 60–69.** Source 68 bounds the principal cost uniformly in the arity m and the order a , with $\mathcal{I}_{m,a}/m \rightarrow 1 + 1/(2c)$ when $a = m^{c+o(1)}$ (its Theorem II.313.1 and Corollary II.313.6), the growing-order regime this theorem leaves open.

For arbitrary targets the relevant worst possible discrepancy has order two. Thus the leading constant is $1/(2 \log 2)$ in the exponent cost of this filtering scheme. Theorem B is a limitation of a specified method, not a lower bound on every possible combinatorial restriction theorem.

F.63.1 Source comparison and contribution boundary

The point of departure is the original paper and the research materials already present in the requested repository. The repository was inspected at commit `5d4718772d1a7dd4c48e1f49c8f92390c73b7ff1`, dated 7 October 2026. Its combined report contains sources 01–27; later source packages and their audits were also consulted where directly relevant. The principal comparisons are as follows.

Source	Existing contribution	Relation to this article
Gowers [1], Lemma 9.3	All-modulus restriction with sixteen-variable exponent 2^{19} .	Theorem A retains the all-modulus scope and gives exponent 54 with a numerical size hypothesis.
Shkredov [235], Lemma 2.1	A broader restriction construction for systems of equations, with large explicit parameter losses.	A relevant predecessor for the selection architecture; our quantitative claims concern one balanced relation.
PROVEIT source 04 [236]	Powered cosine filters, weighted exceptional-relation control, and polynomial restriction over prime fields.	These ingredients are credited as prior work; we amplify a fixed filter and account for small torsion.
PROVEIT source 17 [236]	A single Fejér filter, a signed selection objective, and prime-field retention with powers $\alpha^m \eta^{m-1}$.	Those powers are better in that restricted setting. The present result treats every cyclic modulus.
PROVEIT sources 29 and 40 [236]	Finite-field rank refinements; an exact small-torsion principal spectrum in a one-filter arrangement argument.	The spectrum itself is not claimed as new. The decaying multi-layer error and uniform all-modulus theorem are the contributions here.

In particular, source 04 already proves, for $m = 2r$, a prime-field restriction with retained coefficient

$$\frac{\alpha^{2r} \eta^{2r-1}}{2(400r)^r}$$

under an explicit threshold of power $2r$ in $(\alpha\eta)^{-1}$. Source 17 gives the coefficient $\alpha^m \eta^{m-1}/(2(m+2)^m)$ with a different threshold. Neither should be replaced by a claim that exponent 54 improves all known prime-field estimates. The new statement instead answers the small-prime composite-modulus extension problem recorded in the repository’s research questions, in the form of a uniform cyclic-group theorem and a torsion-sensitive finite-group theorem.

The literature comparison is targeted, rather than a claim of exhaustive priority verification. We have not found the present expectation estimate, the numerical all-modulus corollary, or Theorem B in the inspected sources. The proofs below are self-contained. Modern improvements to the global Szemerédi bound, such as [2], use substantially more of the inverse-theorem machinery; this article makes no claim to improve those global results.

[write] 7 October 2026; sources 54–59. *Identifications and credit within this report.* [236] is this report and its directory at 5d4718772, an ancestor of HEAD (the arrival commit of sources 51–53); the article blob it read (9897851) held sources 01–27, the README blob b2e25fa, and Sections08_09.lean (0fb6e0a) and Definitions.lean (97113b7) are unchanged at HEAD fe7165a3c. Its sources 04, 17, 29 and 40 are this report’s (Sections F.9, F.22, F.34 and F.45); of source 40 it read the audit and verifier, not the manuscript, as it says. The “small-prime composite-modulus extension problem recorded in the repository’s research questions” is source 17’s Question 4 (item II.8). *Not cited, and not citing it:* source 48 (Section II.186), whose constructive restriction also counts repeated entries; both arrived on 6 October 2026. Later members placed after this source and not yet written in (sources 61, 62 and 68) give other restrictions on every cyclic modulus, among them exponent 53 (source 68). Source 58’s three programs were rerun on copies at the write.

F.64 Source 59. Scope, main results, and relation to the source

Source 59, delivered Section 1.

Added in the tenth write, 7 October 2026. Delivered title block. “Sharp Character Separation and Affine Rigidity in Gowers’ Inverse Steps / Torsion-sensitive Freiman restrictions, signed energy inequalities, and a research programme for the ProveIt Ramsey project / Research manuscript prepared with ChatGPT for Vladimir Reshetnikov / 7 October 2026”; *the PDF author field reads* “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov”. The title page is dated 7 October 2026; the archive arrived at 21:44 on 6 October 2026 (Pacific time). Archive gowers-character-affine-refinements.zip (arrival 9458a55a9), manuscript gowers-character-affine-refinements.tex (1,878 lines, 32 pages), pin bec460a33 (the sixth write of this report; its source_manifest.json, shipped as data/59-character-affine-source_manifest.json, records the article, the README and the transcription by blob and SHA-256). Printed below: the abstract, the status statement after it and Section 1; Sections 7–10 open the Part II chapter on affine rigidity (Section II.248), Sections 2–6 form the Part II chapter on character separation at Section 7 after it (Section II.261), Sections 11–12 are in Part V. Its macro \repo (in small capitals) is renamed \reposc, as source 58’s; its ‘question’ environment (“Question”, numbered within sections) is set in this report’s ‘Question’ style.

Abstract of source 59 (as delivered). We develop two local refinements of the additive-combinatorial inverse steps surrounding Sections 6 and 7 of Gowers’s proof of Szemerédi’s theorem. First, a strict averaging argument gives the exact worst-case number of characters needed to separate a finite symmetric forbidden set at explicit radii. If the set has q sign orbits, the universal answer at radius $1/5$ is $\lfloor \log_2(q+1) \rfloor$; for odd-order defects at radius $1/7$ it is $\lfloor \log_3(2q+1) \rfloor$. Both ranks, and the corresponding rank–radius pairs, are sharp. A prime-primary variant, exact order profiles, finite-image rounding, and weighted localization yield Freiman restrictions over arbitrary abelian targets, independently of target size or rank. This refines a classical Green–Tao separation method; an order-five calculation also identifies a failure in a displayed probability estimate in its published proof.

Second, an exact signed graph expansion and a three-position kernel average strengthen the local cubic affine-rigidity profile. For targets without involutions, the sufficient error-support radius increases from the repository’s $1/7$ to $4/9$; excluding order-three elements increases it to $1/2$. Equality requires a single nonzero error value on a subgroup coset; the first endpoint is strict, and the second has precisely the index-two coset extremizers. Quantitative support and color remainders are retained, with sharp support coefficients in the underlying kernel bounds. Classical homomorphism correction then gives sharp-profile recovery for graph defect below $2/9$ and the exact non-affine energy maximum for domains of size at least 16; these global conclusions are credited as applications of prior correction theory. Full proofs, exact verification scripts, source comparisons, and concrete further questions are included. No improved global bound in Szemerédi’s theorem or solution of the lonely runner conjecture is asserted.

Status. This is an AI-assisted research draft with written proofs and independent mathematical and computational checks. It has not been externally refereed or formally verified. Claims of refinement are relative to the cited statements and the pinned repository; mathematical priority remains subject to specialist review.

Gowers’s proof of Szemerédi’s theorem passes from analytic non-uniformity to structured additive relations, and then from many respected relations to an exact local algebraic object. Two interfaces are especially suitable for a precise audit: the graph-energy conclusion of Proposition 6.1 and the extraction of a higher-order Freiman restriction in Lemma 7.5 [1]. This article concentrates on these interfaces. It also develops the near-extremal affine problem that the PROVEIT project places next to Proposition 6.1.

The repository was inspected at commit

bec460a3335f8da421459b4bf69d082e674560a6.

The comparison uses the combined article and research README in `Combinatorics/Ramsey/Research/GowersSzemeredi/local-quantitative-refinements/`. In particular, “source 11”, “source 14”, and “source 24” below mean the internally numbered manuscripts in that pinned project, not separately refereed publications [240]. The combined article contains sources 01–33; the inspected README also records later reports through source 46. This distinction matters when describing what has already been done.

F.64.1 Exact separation of forbidden values

Let H be abelian and let $F \subset H \setminus \{0\}$ be finite and symmetric. Count one representative of each orbit under $f \mapsto -f$, including a singleton for an involution, and write $q = |F/\{\pm 1\}|$. A family of characters $\chi_j : H \rightarrow \mathbb{T}$ separates F at radius θ if every $f \in F$ satisfies $\|\chi_j(f)\|_{\mathbb{T}} \geq \theta$ for at least one j . Here $\mathbb{T} = \mathbb{R}/\mathbb{Z}$ and $\|\cdot\|_{\mathbb{T}}$ is distance from the nearest integer.

Theorem II.263.4 proves the following exact worst-case bounds:

Finite orders allowed among F	Radius	Number of characters
Arbitrary	$1/5$	$\lfloor \log_2(q+1) \rfloor$
Odd	$1/7$	$\lfloor \log_3(2q+1) \rfloor$
Powers of a fixed prime p	$1/(2p)$	$\lfloor \log_p((p-1)q+1) \rfloor$

Infinite-order defects are allowed in every row. The conditions concern the defects, and need not hold for the whole target group. Each rank is sharp for every positive integer q . The radii $1/5$ and $1/7$ are sharp if their respective rank bounds are retained for every q . The prime-primary radius has a different, more limited optimality statement: it is best for its uniform probability contraction, while its optimal separation radius remains undetermined here.

The key gain is strict averaging. At a stage with s forbidden sign orbits, the expected number surviving one character is at most s/b . A neighborhood of the zero character makes all s survive. Thus some character leaves strictly fewer than s/b , so the surviving count is at most $\lfloor (s-1)/b \rfloor$. Iterating this integer recurrence gives the displayed ranks. Projective-line obstructions give matching lower bounds.

For the graph $\Gamma = \{(x, \phi(x)) : x \in B\}$ of a map on a finite subset of an abelian group, let

$$F_k^0 = \{h \in H : (0, h) \in k\Gamma - k\Gamma\}, \quad F_k = F_k^0 \setminus \{0\}. \quad (\text{F.64.1})$$

Localizing all character values to translated short arcs produces a subset on which ϕ is an exact Freiman homomorphism of order k . In the universal case its size is at least

$$\frac{|B|}{(5k)^{\lfloor \log_2(q+1) \rfloor}}, \quad q = |F_k/\{\pm 1\}|. \quad (\text{F.64.2})$$

The same bound holds for any prescribed nonnegative weight in place of cardinality. Finite character images improve the loss by an exact rounding formula. For targets of exponent p , this recovers finite-field fibers with loss $p^{\lfloor \log_p((p-1)q+1) \rfloor}$, independently of k . For infinite-order defects, one character instead gives a linear loss $2kq$.

These are refinements of an established architecture. Green–Tao explicitly use character separation and Bohr localization in Lemmas 8.3 and 9.2 of their U^3 inverse theorem [4]. Gowers–Milićević use a random finite-field kernel to avoid vertical defects in Lemma 3.3 [237]. We do not claim the architecture or dimension-independent finite-field extraction as new. The exact ranks, order accounting, weighted formulation, and finite-image optimization are the refinements developed here.

F.64.2 A stronger affine-rigidity inequality

For $f : G \rightarrow W$, with $|G| = N$, let $\mathcal{E}(f)$ count ordered quadruples respecting addition both in the domain and in the target, and set $\varepsilon(f) = 1 - \mathcal{E}(f)/N^3$. Suppose $W[2] = \{0\}$, where $W[r] = \{w : rw = 0\}$. Subtract an arbitrary affine reference map ℓ and let $S = \{x : f(x) \neq \ell(x)\}$, $t = |S|/N$. If c_b counts the nonzero error value b , define

$$a = \frac{|S|^2 - \sum_{b \neq 0} c_b^2}{N^2}, \quad \Psi(t) = 4t - 10t^2 + 6t^3.$$

Write $d = t^3 - E(S)/N^3$ for the normalized support-energy deficit. Theorem II.249.3 establishes the structural inequality

$$\varepsilon(f) \geq \Psi(t) + 6 \left(t^3 - \frac{E(S)}{N^3} \right) + 2(1 - 3t)a, \quad 0 \leq t \leq \frac{1}{2}. \quad (\text{F.64.3})$$

It improves the repository’s color coefficient $2(1 - 7t)$ and, below $t = 1/3$, separately controls failure of the support to be a coset and failure of the errors to have one value.

A second estimate averages the signed quartic kernel over the three ways to choose one negative position among three free positions. It gives the larger sufficient ranges in Theorem II.250.4:

$$\varepsilon(f) \geq \Psi(t) + 3d + \left(2 - \frac{9}{2}t \right) a, \quad 0 < t < \frac{4}{9}, \quad W[2] = \{0\}; \quad (\text{F.64.4})$$

$$\varepsilon(f) \geq \Psi(t) + 4d + \min\{2 - 3t, 4 - 8t\} a, \quad 0 < t < \frac{1}{2}, \quad W[2] = W[3] = \{0\}. \quad (\text{F.64.5})$$

In either open range, equality with $\Psi(t)$ holds exactly when the error is one nonzero constant on a subgroup coset. These are statements relative to any affine reference; no nearest-map hypothesis is needed. The radii are sufficient, not asserted optimal.

Corollary II.250.6 also resolves both boundaries: at $t = 4/9$ the first class has strict inequality, while at $t = 1/2$ in the second class equality holds exactly for a single nonzero error value on an index-two coset. The coefficients 3 and 4 are sharp in the underlying signed-kernel bounds with their other terms fixed.

The cubic itself and its coset extremizers are already in source 11. Our contribution is a stronger support/color decomposition and an enlarged local range. In contrast, the global recovery range $\varepsilon < 2/9$, the cutoff $N \geq 16$ for the exact non-affine maximum, and the corresponding partial-graph application can already be obtained by combining that prior cubic with the published correction theorem of Ben-Or–Coppersmith–Luby–Rubinfeld [27, 28]. We supply complete proofs and record the improved repository consequences, without presenting them as a new homomorphism-testing theorem.

F.64.3 What the bounds do and do not improve

The results are local improvements to the inverse step. If $|\Gamma - \Gamma| \leq C|\Gamma|$, then $|F_k^0| \leq C^{2k+1}$ by a classical sum-difference estimate and a vertical injection. Thus (F.64.2) is polynomial in C , uniformly over arbitrary abelian targets. This answers the concrete request for an explicit mixed-torsion continuation of the repository’s finite-field restriction, though it does not optimize its exponent and does not improve the best prime-cyclic bound already in source 10.

For comparison, Gowers’s original Lemma 7.5 retains $|B|/(8kC^{4k})$ in its prime-cyclic setting. The present universal exponent $(2k+1)\log_2(5k)$ can be larger than $4k$. Its advantage is its unrestricted target, exact forbidden-set parameter, and torsion-sensitive alternatives. Any comparison should keep all those hypotheses visible.

An improvement to the final bound for progression-free sets would require compatible gains through the remaining inverse, phase-localization, and density-increment steps. We do not provide that chain. Modern global bounds, such as those of Leng–Sah–Sawhney [2], involve substantially more than the interfaces studied here. Similarly, the one-character problem contains the integer lonely runner conjecture exactly; we identify that connection and transfer known bounds, rather than claiming to solve it.

F.64.4 Conventions

All groups are written additively, and $C_n = \mathbb{Z}/n\mathbb{Z}$. A character means a homomorphism into $\mathbb{T}$, with the domain given the discrete topology. For finite sets A , the notation $rA - sA$ allows repetitions. Energies are unnormalized ordered counts. A Freiman homomorphism of order k respects every equality of two k -term sums, including repeated summands. A finite nonempty set is a coset precisely when it is closed under $x - y + z$. We distinguish q , the number of sign orbits, from $|F|$: if $i(F)$ counts the nonzero involutions in F , then

$$q = \frac{|F| + i(F)}{2}.$$

All probability arguments can be restricted to a finitely generated subgroup containing the finitely many values under discussion. This keeps infinite-target assertions separate from any finiteness assumption on the domain.

[write] **7 October 2026; sources 54–59.** *Identifications within this report.* [240] is this report at bec460a33, its sixth write (sources 01–33; the blobs and SHA-256 values recorded in `data/59-character-affine-source_manifest.json` were verified against that commit), and its “source 10”, “source 11”, “source 14”, “source 24” and “source 31” are this report’s (Sections F.15, F.16, F.19, F.29 and F.36). The repository’s “1/7” and colour coefficient $2(1 - 7t)$ are source 11’s Theorems F.16.1 and II.25.3; its “sufficient cutoff 34” is source 11’s Theorem F.16.3, with the third write’s remark on smaller groups (Remark W2); the “explicitly unfinished correction/cutoff step” is item II.21; the mixed-torsion request is source 24’s Question 7

(item II.8); and the support-coset stability already answered by combining sources 11 and 14 is item II.20. The suite’s three `make check` programs were rerun on a copy at the write.

F.65 Source 60. Research context and principal conclusions

Source 60, delivered Section 1.

Added in the eleventh write, 7 October 2026. Delivered title block. “Fourier limits, exact arrangement repair, and translated restriction in Gowers’s quantitative framework / Research prepared with ChatGPT for Vladimir Reshetnikov / ProveIt research programme / 7 October 2026 | Research manuscript, version 1.0”; *the PDF author field reads “Research prepared with ChatGPT for Vladimir Reshetnikov”*. The title page is dated 7 October 2026; the archive arrived at 22:07 on 6 October 2026 (Pacific time). Archive `Gowers_Fourier_Arrangement_Refinements.zip` (arrival `0d93c1eac`), manuscript `gowers_fourier_arrangements.tex` (3,145 lines, 49 pages), pin `89c1cbc38` (“7 October 2026, 04:12:28 UTC”; its `audit/SOURCE_AUDIT.json`, shipped as `data/60-fourier-limits-SOURCE_AUDIT.json`, records the README, this article and the transcription by blob). Printed below: the abstract, its scope paragraph and Section 1; Section 2 (the norm limit) stands in Part III after source 63 (Section III.94), Sections 3–8 form a chapter of Part II on arrangement repair (Section II.331), Section 9 (translated restriction) stands in the Part II restriction chapter (Section II.310), Sections 10–12 and Appendix A are in Part V. Its macros `\V` (calligraphic $\mathcal{V}$, not this report’s V_d) and `\pin` (its own pin) are renamed `\Vcal` and `\pinsixty`; its ‘problem’ environment (“Research question”, numbered 1–14 on its own counter) is set in this report’s ‘Research question’ style, numbered within its section, which keeps the indices; its figure is shipped as `figures/60-fourier-limits-arrangement_profiles.pdf`. It loads `cleveref` and does not use it.

Abstract of source 60 (as delivered). We develop four quantitative refinements around Gowers’s proof of Szemerédi’s theorem and the research already collected in the ProveIt repository. First, the limiting comparison constant between the fourth powers of the real, centered U^2 and U^d norms on finite odd-order abelian groups is exactly a bounded Fourier extremum. A uniform Fourier-tail estimate justifies passage through varying groups and unbounded functions. It reduces the previous interval $[1/3, 1/2]$ to

$$\frac{1}{3} \leq c_{\infty}^{\text{odd}} \leq \frac{1}{2} - \frac{8}{\pi^2} + \frac{64}{\pi^4} = 0.3464533951\dots$$

When the degree and the least prime divisor of the group order both tend to infinity, the exact limiting constant is $1/3$, by combining the compactness theorem with a Riesz–Sobolev inequality of Christ and Iliopoulou.

Second, for every even number $s \geq 4$ of vertical differences in a respected-arrangement test, we prove an explicit polynomial lower bound for its defect, classify every interior equality case, and determine every coefficient of the optimal small-error inverse expansion. This resolves the higher-order coefficient and exact-equality questions left by repository source 33. Third, we prove a sharp robust-intersection theorem for the two directional model classes, and obtain biaffine repair with optimal inverse coefficients through degree $s - 2$. Finally, a translated Fejér selector retains the existing polynomial restriction count while reducing the prime-uniform size threshold from a proposed power L^m to L^{m-1} ; an exact certificate records small torsion and repeated tuples. All new assertions receive proofs. Exact finite diagnostics, source provenance, and a prioritized research programme accompany the article.

Scope of the claims. The advances are local extremal and restriction theorems. They do not supply a new global bound for progression-free sets, close Gowers’s entire quantitative induction, or constitute a Lean formalization. Claims of improvement are made against the inspected repository snapshot, with published inputs identified explicitly; no claim of publication priority is inferred from a literature search.

Gowers’s proof of Szemerédi’s theorem links uniformity norms, additive relations on graphs, restriction to approximate homomorphisms, and structural repair [1]. Several estimates in this chain have a meaningful extremal problem of their own. Improving such an estimate can clarify which configurations force a loss, even when a global density-increment iteration has additional bottlenecks.

The present article studies four such problems. They are connected by a common principle: retain the distribution or Fourier mass responsible for an error before replacing it by a worst-case bound. For high-order norm comparisons this means controlling the whole tail of the Fourier fourth powers. For arrangement repair it means coupling the constant Fourier coefficient with the remaining spectrum. For two-direction testing it means retaining both conditional variances. For restriction it means summing the actual nonprincipal Fourier weights after an additional translation average.

F.65.1 The baseline and the research standard

The repository was inspected at commit

89c1cbc38600279fed03c8e741aba9eafb08ca70,

dated 7 October 2026, 04:12:28 UTC. At that snapshot the consolidated research article incorporates sources 01–33, while sources 34–46 are separately staged. We compared the main targets with the consolidated article, its source inventory and open questions, and the relevant later audits. The numbered “repository sources” below refer to this inventory, not to published papers [243].

The closest predecessors are source 29 for odd-order norm comparisons, source 33 for high-agreement arrangement repair, and sources 04 and 17 for polynomial restriction [247, 248, 244, 245]. Source 45 records a stronger finite- U^4 comparison than source 29; our asymptotic theorem does not replace that finite-degree estimate. A crosswalk and exact source identifiers are supplied in Section V.224 and in the archive.

Every theorem asserted here has a written proof, including its normalization, range of parameters, and use of earlier results. Separate derivations were compared during preparation, and bounded exact computations were used to test equality cases, small torsion, repeated coordinates, and algebraic constants. These checks are evidence against implementation and algebra errors. They are not substitutes for the universal proofs and have not been certified by a proof assistant.

F.65.2 A limiting norm problem becomes a bounded Fourier problem

For $d \geq 2$, define

$$c_d^{\text{odd}} = \sup_{G,f} \frac{\|f\|_{U^2(G)}^4}{\|f\|_{U^d(G)}^4}, \quad \beta_{\text{odd}} = \sup_{G,f} \frac{\|f\|_{U^2(G)}^4}{\|f\|_{\infty}^4}, \quad (\text{F.65.1})$$

where G ranges over finite abelian groups of odd order and $f : G \rightarrow \mathbb{R}$ ranges over nonzero functions with $\mathbb{E}f = 0$. All averages are normalized. Monotonicity gives a decreasing sequence c_d^{odd} . Source 29 proves that its limit lies in $[1/3, 1/2]$ [247].

We prove that

$$\lim_{d \rightarrow \infty} c_d^{\text{odd}} = \beta_{\text{odd}}, \quad \frac{1}{3} \leq \beta_{\text{odd}} \leq \frac{1}{2} - \frac{8}{\pi^2} + \frac{64}{\pi^4}. \quad (\text{F.65.2})$$

For a fixed finite group, $\|f\|_{U^d} \rightarrow \|f\|_{\infty}$ is not enough to prove this assertion: the group and the function in the supremum can change with d . The central new estimate prevents Fourier

fourth-power mass from escaping to more and more small coefficients. A compact dual-group construction then preserves the relevant mass while converting high-degree norm control into an L^∞ bound.

The final numerical upper bound uses the sharp coefficient inequality

$$|\widehat{f}(\chi)| \leq \frac{1}{q} \cot \frac{\pi}{2q} < \frac{2}{\pi}$$

for real centered $|f| \leq 1$ and a character of odd order q . Parseval then constrains how the remaining Fourier mass can be distributed. The resulting interval has width approximately 0.01312006183, compared with $1/6$ previously. In particular, the possibility that the limit is $1/2$ is excluded.

There is also an exact result in a more specific limiting regime. If the least prime divisor of $|G|$ tends to infinity along with d , the compact limit is connected. The Riesz–Sobolev inequality in [242, Theorem 1.2] then gives the exact limiting value $1/3$. Neither conclusion supplies an effective rate in d . The unrestricted identity $\beta_{\text{odd}} = 1/3$ remains a precise open problem.

F.65.3 The complete arrangement polynomial and its extremizers

Let G, H, K be finite abelian groups and $s \geq 4$ even. Sample $h \in H$, independent $y_1, \dots, y_s \in H$, and a uniform balanced relation

$$x_1 + \dots + x_{s/2} = x_{s/2+1} + \dots + x_s$$

in G . A map $\phi : G \times H \rightarrow K$ passes when the same balanced relation holds among its vertical differences $D_h \phi(x_i, y_i)$. Write $\varepsilon_s(\phi)$ for the failure probability. Its exact model class is

$$\mathcal{V} = \{a(x) + B(x, y) + \ell(y) : a \text{ arbitrary, } B \text{ biadditive, } \ell \text{ additive}\}.$$

Let δ be the normalized Hamming distance to $\mathcal{V}$.

The new polynomial is

$$B_s(\delta) = \frac{1 - (1 - 4\delta)^s}{4} - 2^{2s-3} \delta^{s-1} (1 - 2\delta). \quad (\text{F.65.3})$$

For $0 \leq \delta \leq 1/4$ we prove $\varepsilon_s \geq B_s(\delta)$. For every even $s \geq 6$, this strictly improves the earlier polynomial throughout $0 < \delta < 1/4$. At $s = 4$ it recovers the earlier sharp cubic $4\delta - 24\delta^2 + 32\delta^3$ and adds its complete interior equality classification.

For $0 < \delta < 1/4$, equality holds exactly when, after adding an exact model,

$$\phi(x, y) = a \mathbf{1}_C(x) \ell_2(y), \quad 2a = 0, \quad a \neq 0, \quad (\text{F.65.4})$$

where C is a coset of a subgroup of G of index $n \geq 3$, and $\ell_2 : H \rightarrow \mathbb{Z}/2\mathbb{Z}$ is a surjective homomorphism. Then $\delta = 1/(2n)$. Both even and odd indices are allowed. Thus the extremal errors fill half of the points in a coset of horizontal rows.

Under the inherited explicit activation condition $\varepsilon_s \leq 1/[240(s-1)]$, the appropriate small inverse branch $\rho_s = B_s^{-1}$ satisfies $\delta \leq \rho_s(\varepsilon_s)$. For $s \geq 6$,

$$\rho_s(u) = \frac{u}{s} + \frac{2(s-1)}{s^2} u^2 + \frac{8(s-1)(2s-1)}{3s^3} u^3 + O_s(u^4). \quad (\text{F.65.5})$$

Every Taylor coefficient is optimal in succession, with all preceding coefficients fixed. This is stronger than identifying the cubic coefficient alone. It does not mean that B_s is the exact pointwise envelope at every real distance: its nonzero equality distances are discrete.

For Gowers's eight-arrangements, $s = 16$, so

$$\delta \leq \frac{\varepsilon}{16} + \frac{15}{128}\varepsilon^2 + \frac{155}{512}\varepsilon^3 + O(\varepsilon^4).$$

The common-prime case, in which all three groups equal $\mathbb{F}_p$, has a different second-order extremal problem. A separate construction and an odd-target cancellation estimate improve both sides of its coefficient bracket:

$$\frac{s-1}{s^2} \leq c_s^{\text{prime}} \leq \frac{3s-2}{2s^2}.$$

For $s = 16$ this is $15/256 \leq c_{16}^{\text{prime}} \leq 23/256$. It does not determine the optimal coefficient.

F.65.4 Removing the arbitrary row function by a second test

Let $\mathcal{H}$ be the horizontal analogue of $\mathcal{V}$, and let

$$\mathcal{A} = \{c + B(x, y) + m(x) + \ell(y) : B \text{ biadditive, } m, \ell \text{ additive}\}$$

be the fully biaffine class. If $d_v = \text{dist}(\phi, \mathcal{V})$, $d_h = \text{dist}(\phi, \mathcal{H})$ and $\eta = d_v + d_h < 1/4$, we prove

$$\text{dist}(\phi, \mathcal{A}) \leq J(\eta) := \frac{1 + \eta - \sqrt{1 - 2\eta}}{2}. \quad (\text{F.65.6})$$

Binary unions of one horizontal and one vertical strip show that this is the optimal continuous modulus from η alone.

Write $e = \varepsilon_s^v + \varepsilon_s^h$ for the sum of directional failure probabilities and $\delta_A = \text{dist}(\phi, \mathcal{A})$. The local defect bound is

$$e \geq P_s(\delta_A) := 2(1 - \delta_A)B_s(\delta_A) \quad (\delta_A \leq 1/4). \quad (\text{F.65.7})$$

It gives an explicit inverse repair theorem when $e \leq 1/(1920s)$, with

$$\delta_A \leq \frac{e}{2s} + \frac{2s-1}{4s^2}e^2 + O_s(e^3).$$

More generally, all inverse coefficients through degree $s-2$ are optimal in succession. This answers source 33's two-direction question locally, including its request for sharp first-order constants. The full finite-error envelope is still open.

F.65.5 A stronger restriction threshold with the same retention

For even $m \geq 4$, let $B \subseteq \mathbb{F}_p$ have size b , and suppose a map $\phi : B \rightarrow \mathbb{F}_p$ respects at least αb^{m-1} ordered balanced m -tuples. For $0 < \alpha, \eta \leq 1$, put

$$L = \left\lceil \frac{m+1}{\alpha\eta} \right\rceil.$$

We prove that the condition

$$b \geq 2 \left(1 + \binom{m}{2} \right) L^{m-1} \quad (\text{F.65.8})$$

guarantees a subset $B' \subseteq B$ on which at least a fraction $1 - \eta$ of the balanced tuples are respected and whose total number of balanced tuples is at least

$$\frac{\alpha^m \eta^{m-1}}{2(m+2)^m} b^{m-1}. \quad (\text{F.65.9})$$

This retains source 17’s coefficient while improving the proposed combined threshold of sources 04 and 17 by one power of L . It also improves source 17’s original threshold, whose degree is $2m - 1$ in the inverse accuracy parameters. Fixed-characteristic affine-rank restriction is a separate regime where source 29 has stronger parameter dependence; our assertion is uniform as the prime modulus varies.

At $m = 16$, a respected-quadruple count γb^3 yields a respected 16-tuple count $\gamma^7 b^{15}$ by graph-energy interpolation. Thus one may take

$$L = \left\lceil \frac{17}{\gamma^7 \eta} \right\rceil, \quad b \geq 242L^{15},$$

and retain at least $\gamma^{112} \eta^{15} b^{15} / (2 \cdot 18^{16})$ balanced 16-tuples. The proof also gives an exact finite-group certificate that records short-torsion leakage instead of assuming it absent.

F.65.6 Notation and organization

Every finite set carries its uniform probability measure. Compact abelian groups carry normalized Haar measure. Unless a section explicitly specifies a different local convention, character groups are written additively when recording algebraic relations and multiplicatively when evaluating characters; $\bar{\chi} = \chi^{-1}$ makes the convention explicit. Fourier coefficients are normalized, but frequency sums are not:

$$\hat{f}(\chi) = \mathbb{E}_x f(x) \overline{\chi(x)}, \quad \sum_{\chi} |\hat{f}(\chi)|^2 = \mathbb{E}_x |f(x)|^2.$$

For maps with values in a finite group, distance always means the uniform fraction of unequal values. The symbol s denotes the even number of vertical differences in an arrangement; an s -difference arrangement has $2s$ point occurrences. In the restriction section m denotes the even number of entries in a balanced relation. Elsewhere m can be a local moment index or the half-arity $s/2$, as specified there. The restriction section uses additive circle-valued characters and explicitly switches to unnormalized Fourier coefficients in its graph-energy calculation.

The proofs of the norm limits, arrangement bounds, two-direction repair and restriction theorem occupy the next sections. The final sections discuss integration, computational evidence, and further research. The coarse majority decoder needed only to activate the local inverse bounds is reproduced in an appendix and credited to its repository and testing predecessors.

[write] 7 October 2026; sources 60–69. *Identifications within this report.* [243] is this repository at 89c1cbc38, an ancestor of HEAD; the README (blob 689dbd3), this article (32087ba, the sixth write, sources 01–33) and the transcription (41085a0) recorded in its source audit (shipped as `data/60-fourier-limits-SOURCE_AUDIT.json`) were verified against that commit. Its “repository sources” 04, 17, 19, 29 and 33 are this report’s (Sections F.9, F.22, F.24, F.34, F.38), and its “separately staged” source 45 is this report’s (Section F.50); the finite-degree bound it quotes from source 45, $4795/832$, has since been improved by source 49 to $29/4$ (item III.1). *Since its arrival:* source 63, which arrived in the same commit, proves the bounded inequality that settles this source’s norm limit (Section III.94). The four verifiers of this source were rerun on copies at the write. As this source suggests, it is registered as the next source, number 60, and its labels are kept behind the prefixes `gsr:fm:60:`, `gsr:cc:60:`, `gsr:en:60:` and `gsr:fz:60:`.

F.66 Source 61. Research objective, source comparison, and main results

Source 61, delivered Section 1.

Added in the eleventh write, 7 October 2026. Delivered title page. “Combinatorics / Ramsey Theory / Exact Freiman Restriction in Arbitrary Finite Abelian Groups / Torsion-uniform coloring, relation-energy retention, and logarithmic order barriers / Mathematical development and exact computational checks with ChatGPT / Prepared for the ProveIt project / 6 October 2026”; *the PDF author field reads “ChatGPT; prepared for the ProveIt project”*. Archive `Gowers_Torsion_Exact_Restriction.zip` (arrival `0d93c1eac`), manuscript `ramsey_torsion_exact_restriction.tex` (1,249 lines, 23 pages), pin `5ea1d06dd` (its `SOURCE_AUDIT.md`, shipped as `61-exact-restriction-SOURCE_AUDIT.md`, records three blobs). Printed below: the abstract, the status statement of its title page and Section 1; Sections 2–9 stand in the Part II restriction chapter (Section II.310), Sections 10–13 and Appendix A are in Part V. Its macro `\repo` (in small capitals) is renamed `\repsc`, as for sources 58 and 59; its ‘question’ environment (“Question”, on the shared counter) is set in this report’s ‘Question’ style; its ‘conjecture’ environment is defined and not used.

Abstract of source 61 (as delivered). We develop an exact-restriction alternative to two local steps in Gowers’s proof of Szemerédi’s theorem. If the graph Γ of a map between finite abelian groups satisfies $|\Gamma - \Gamma| \leq K|\Gamma|$, then its entire domain can be partitioned into at most

$$5mK^{(2m+1)\log_2(5m)}$$

sets on which the map is a Freiman homomorphism of order m . No assumption on the orders, ranks, or exponents of the ambient groups is imposed. The partition retains weighted additive energies and any one prescribed scalar correlation with explicit losses.

Combining this with a self-contained, deliberately conservative Balog–Szemerédi–Gowers argument gives an order-eight restriction containing at least $2^{-2193}\alpha^{911}|B|$ points and at least $2^{-35277}\alpha^{14655}|B|^{15}$ additive sixteen-tuples, all respected exactly, whenever the original graph has at least $\alpha|B|^{15}$ such tuples. This strengthens the numerical conclusion of Gowers’s Lemma 9.3, removes its size threshold, and replaces approximate respect by exact respect.

The underlying forbidden-difference coloring problem has a universal power exponent b_m satisfying

$$\lfloor \log_2 m \rfloor + 1 \leq b_m \leq \log_2(5m) \quad (m \geq 2).$$

Thus its leading logarithmic dependence on the relation order is determined. Explicit 2-primary examples also force every exact graph restriction to be a singleton and quantify a separate obstruction to improving the graph theorem. Complete proofs, exact finite checks, an integration plan, and further research problems are supplied.

Status and scope. This is an AI-assisted, unrefereed research manuscript. The stated results are supported by the proofs below; their formulations are offered for mathematical evaluation, not certified publication priority. No proof-assistant verification or improved global Szemerédi bound is claimed. Existing lemmas used as ingredients are explicitly attributed.

The restriction of an approximately additive map to an exactly additive domain is one of the structural transitions in Gowers’s argument [1]. Two quantities matter simultaneously: how much of the domain survives, and how many additive relations survive. A large subset by itself does not give the desired relation count at the original density normalization. Moreover, an argument that chooses a scalar multiplier in a prime field does not automatically extend to a group with mixed torsion.

This manuscript separates the restriction mechanism into three steps. First, a finite forbidden set in the codomain is separated from zero by a short list of characters. Second, narrow character cells prevent every balanced relation of bounded length from having a forbidden defect. Third, graph small-difference estimates bound the number of possible vertical defects. This produces a partition rather than a single favorable restriction, which makes several retention statements automatic.

F.66.1 What is and is not being improved

The repository snapshot used here is

5ea1d06dd2f2d3a30e96016b236e3f1e3dfb048f.

The formal statement of Lemma 7.5 in that snapshot assumes a prime modulus and supplies a restriction of size $|B|/(8mK^{4m})$. Lemma 9.3 is stated for all nonzero cyclic moduli, with a sufficiently-large-modulus hypothesis and an approximate-respect conclusion. The file `Proofs09Restriction.lean` already contains a theorem named `lemma_9_3_holds`. Thus this manuscript does *not* announce that an unproved repository statement has been closed. It gives a stronger, different route to the mathematical conclusion [250, 251].

The research collection already records substantially better *prime-modulus* restrictions, including an order-eight size bound $2^{-209}\gamma^{103}|B|$ under an additive-energy hypothesis, and polynomial approximate-restriction bounds obtained by other methods [252]. The exponent 911 proved here is not competitive with 103 in that special setting. The contribution advanced for evaluation is instead the arbitrary-finite-abelian-group partition, its relation-energy consequences, and the logarithmic torsion obstruction. Neither the general Balog–Szemerédi–Gowers principle nor the Plünnecke–Ruzsa inequalities are claimed as new.

There is a second important limitation. Keeping a chosen scalar correlation is not the same as preserving the complete derivative spectrum or the compatibility of a family of cross-sections. The latter issues arise later in Gowers’s proof. No step below justifies replacing all of Sections 12–16, and no revised global bound for Szemerédi’s theorem is deduced.

F.66.2 Principal statements

All groups in the principal results are finite and abelian, written additively. For a nonempty set $B \subseteq G$ and a map $\phi : B \rightarrow H$, write

$$\Gamma = \{(x, \phi(x)) : x \in B\}, \quad n = |B|.$$

Definitions of the relation energies appear in Section II.322.

Theorem F.66.1 (Exact partition from graph growth). *Let $m \geq 1$ and $|\Gamma - \Gamma| \leq Kn$. Put*

$$a_m = (2m + 1) \log_2(5m), \quad \mathcal{D}_m(\Gamma) = \{h \in H : (0, h) \in m\Gamma - m\Gamma\}.$$

There is a partition of B into R nonempty sets $B_1, \dots, B_R$ such that every $\phi|_{B_j}$ is a Freiman homomorphism of order m , and

$$R \leq (5m)^{\lceil \log_2 |\mathcal{D}_m(\Gamma)| \rceil} \leq 5mK^{a_m}. \tag{F.66.1}$$

In particular, some exact restriction has size at least $n/(5mK^{a_m})$.

[write] **7 October 2026; sources 60–69.** *Against source 59* (Section II.261; checked at the write). Source 59’s Corollary II.265.2 gives $|B'| \geq C^{-(2k+1)\log_2(5k)}|B|$ from the same hypothesis, larger than this theorem’s $n/(5mK^{a_m})$ by the factor $5m$, and its separation rank $\lfloor \log_2(q+1) \rfloor$ for q sign orbits is at most this theorem’s $\lceil \log_2 |\mathcal{D}_m(\Gamma)| \rceil$; this theorem’s partition of the whole domain, with every cell exact, is its own form.

The defect-sensitive first bound is often preferable to its uniform relaxation. If the defect set consists only of zero, the partition has one cell and nothing is lost.

Theorem F.66.2 (Order-eight relation restriction). *Suppose $0 < \alpha \leq 1$ and*

$$\mathcal{E}_8^\phi(B) \geq \alpha n^{15}.$$

There exists a nonempty $B' \subseteq B$ on which ϕ is an exact Freiman homomorphism of order eight, such that simultaneously

$$|B'| \geq 2^{-2193} \alpha^{911} n, \tag{F.66.2}$$

$$\mathcal{E}_8(B') = \mathcal{E}_8^\phi(B') \geq 2^{-35277} \alpha^{14655} n^{15}. \tag{F.66.3}$$

The same conclusions hold if the hypothesis is replaced by $\mathcal{E}_2^\phi(B) \geq \alpha n^3$.

The constants are intentionally rounded to integral powers of two. A general-order formula, with less rounding and a modular interface for stronger energy-to-growth inputs, is proved in Section II.327.

Theorem F.66.3 (Leading order of the universal coloring exponent). *For $m \geq 2$, let b_m be the infimum of the real numbers b for which a constant C_m exists with the following property: for every finite abelian H and every symmetric $D \subseteq H \setminus \{0\}$, one can partition H into at most $C_m(|D| + 1)^b$ cells S satisfying*

$$(mS - mS) \cap D = \emptyset.$$

Then

$$\lfloor \log_2 m \rfloor + 1 \leq b_m \leq \log_2(5m). \tag{F.66.4}$$

Consequently $b_m = \log_2 m + O(1)$ and $\lim_{m \rightarrow \infty} b_m / \log_2 m = 1$.

The lower bound concerns the coloring problem itself, not merely a particular character construction. It does not prove that the exponent a_m in Theorem F.66.1 is optimal. Section II.329 distinguishes these two optimization problems.

F.66.3 Proof architecture and provenance

Sections II.323–II.325 prove the character-separation and exact-partition statements. The Plünnecke argument follows the minimal-growth method of Petridis [36], with its proof included. Section II.326 supplies a conservative energy-to-growth estimate by a standard popular-difference and dependent-random-choice method; see Gowers [1] and Fox–Sudakov [11] for the surrounding methodology. Section II.327 tracks every constant in the combination. Section II.329 gives the obstructions and proves Theorem F.66.3.

Our comparison with existing work is targeted: the original paper, relevant repository statements and research descriptions, and primary references for the standard ingredients. It is not an exhaustive review of every formulation of Freiman restriction or higher-order difference coloring. The proofs and the source-comparison claims should be evaluated separately from any future claim of priority.

[write] 7 October 2026; sources 60–69. *Identifications within this report.* [250], [251] and [252] are this repository at 5ea1d06dd, an ancestor of HEAD; Sections06_07.lean, Sections08_09.lean and Proofs09Restriction.lean have the blobs recorded in its source audit at HEAD fb9602e55, where lemma_7_5 (line 183) still assumes a prime modulus with restriction size $|B|/(8kC^{4k})$ and lemma_9_3_holds is at line 376 of Proofs09Restriction.lean. The “order-eight size bound $2^{-209}\gamma^{103}|B|$ ” recorded in the research collection is source 10’s Corollary II.21.2 (prime moduli). *Not cited, and not citing it:* source 59, whose character-separation exponent $(2m+1)\log_2(5m)$ this source finds independently (note after Theorem F.66.1); source 59 was placed after this source’s pin. Source 61’s checks were rerun on a copy at the write.

F.67 Source 62. The target and the contribution

Source 62, delivered Section 1.

Added in the eleventh write, 7 October 2026. Delivered title block. “Unit-Dilation Restriction on Cyclic Groups / Polynomial losses for Gowers’ approximate homomorphisms and exact weighted degeneracy constants / Research manuscript prepared with ChatGPT / For the ProveIt Ramsey research collection / 6 October 2026”; *the PDF author field reads “Research manuscript prepared with ChatGPT”.* Archive ProveIt_All_Moduli_Restriction.zip (arrival 0d93c1eac), manuscript article.tex (1,317 lines, 20 pages), pin 17123532a (its SOURCE_AUDIT.md, shipped as 62-unit-dilation-SOURCE_AUDIT.md, records four blobs; data/62-unit-dilation-CLAIMS.json the same revision). Printed below: the abstract, its status statement and Section 1; Sections 2–9 stand in the Part II restriction chapter (Section II.310), Sections 10–13 are in Part V. Its macros \ee ($\exp(2\pi i \cdot)$, not this report’s $e(\cdot)$), \U (roman U, not this report’s $\mathcal{U}$) and \R (calligraphic $\mathcal{R}$, not the reals) are renamed \eexp, \Urm and \Rcal; its ‘question’ environment (“Question”, on the shared counter) is set in this report’s ‘Question’ style.

Abstract of source 62 (as delivered). We give a one-stage restriction argument for maps between cyclic groups of arbitrary, possibly different, orders. The argument extends the polynomial-loss Fejér restriction method beyond prime moduli. Its two main ingredients are averaging the codomain dilation over units, which distinguishes nonzero defects of small order, and a weighted gcd count, which avoids deleting every tuple satisfying an additional short relation.

For the order-eight interface of Gowers’ Lemma 9.3, suppose that $B \subseteq \mathbb{Z}/N\mathbb{Z}$ has size βN and a map $\phi : B \rightarrow \mathbb{Z}/M\mathbb{Z}$ respects at least $\alpha|B|^{15}$ additive sixteen-tuples. A completely elementary version produces a restriction that respects a proportion at least $1 - \eta$, retaining at least $2^{-245}\alpha^{31}\eta^{30}|B|^{15}$ respected tuples whenever $N \geq 2^{253}(\alpha\eta)^{-31}\beta^{-15}$. A more precise parameter choice gives retention $\gg \alpha^{16}\eta^{15}|B|^{15}/[\log \log(e^e/(\alpha\eta))]^{15}$, with a polynomial size threshold; this sharper envelope uses the classical explicit totient estimate of Rosser and Schoenfeld.

We also determine the exact weighted exceptional-tuple constant at each filter length, identify moduli attaining it, and prove that its limit is $\zeta(m-2)/\zeta(m-1)$ for m -tuples. A primordial construction shows that a uniform constant bound for the primitive-dilation response of these kernels is impossible: the double logarithm has the correct worst-case scale for this specific detector. All new statements are proved below. Exact finite checks, an integration crosswalk, and a research agenda accompany the manuscript.

Status and scope. This is an AI-assisted mathematical research draft, not an independently refereed paper or a Lean-kernel-checked development. The claims of proof below refer to the written arguments. No publication-priority claim is made. The result improves a local restriction step; it does not establish a new end-to-end bound for Szemerédi’s theorem or a Ramsey number.

Gowers’ restriction step asks how to turn a positive density of respected additive relations into an almost perfectly respected family on a subset. The method is useful because it does not first

require an exact Freiman homomorphism. Lemma 9.3 of [1] is the order-eight instance. The present article addresses the quantitative restriction problem itself, rather than the subsequent extension or phase-integration steps.

The repository’s current research summary records polynomial-loss restriction results over prime moduli and fixed finite fields, while distinguishing these from the all-moduli interface [254]. The obstruction to simply reusing prime-modulus averaging is real. If a nonzero target defect has order q , averaging a dilation over *all* scalars leaves every Fourier mode whose index is divisible by q . Increasing the filter length does not remove this effect. We calculate the obstruction explicitly in Section II.320.

Our replacement averages the target dilation over units. The resulting response is a Ramanujan-sum transform of a positive trigonometric kernel. This has two advantages: every nonzero defect is tested through primitive q th roots, and the response is bounded independently of the ambient modulus. Separately, we sum the weights of exceptional short relations instead of counting their supports. The resulting error is $O_m(N^{m-2})$, with a constant independent of filter length, even over composite cyclic groups.

The distinction from the existing formal development is important. At the inspected repository revision, the module `Proofs09Restriction.lean` already contains `lemma_9_3_holds` [256]. Thus the contribution is a stronger estimate and a different proof mechanism, not a claim that the old all-moduli statement has no proof.

Main mathematical outputs. The principal restriction theorem is Theorem II.314.1. The fully explicit polynomial order-eight consequence is Corollary II.319.2. The exact arithmetic error formula and its limiting constant are Theorem II.317.1. Theorem II.320.2 identifies a genuine limitation of the unit-Fejér detector. Theorem II.318.3 preserves the entire defect-order distribution, and Theorem II.321.1 extends the argument to labelled domains of bounded fibre multiplicity.

[write] **7 October 2026; sources 60–69.** *Identifications within this report.* [254] is this directory’s README at 17123532a (blob 689dbd3, the sixth write); [255] and [256] are `Sections08_09.lean` and `Proofs09Restriction.lean`, whose recorded blobs are unchanged at HEAD fb9602e55 (`lemma_9_3` at line 77, `lemma_9_3_holds` at line 376). Its “polynomial-loss restriction results over prime moduli and fixed finite fields” are sources 04, 17 and 29 (Sections II.57 and II.119). *Not cited:* source 58, placed after this source’s pin, and source 68, later still; their all-cyclic theorems are compared with this source’s in Section II.310. The suite was rerun on a copy at the write.

F.68 Source 63. Introduction and principal results

Source 63, delivered Section 1.

Added in the eleventh write, 7 October 2026. Delivered title page. “ProveIt / Combinatorics / Ramsey / Asymptotically Sharp Constants in Gowers’ Szemerédi Framework / Odd-order norm limits and efficient local phase removal / Prepared for Vladimir Reshetnikov and the ProveIt project / Mathematical development and manuscript by ChatGPT / 7 October 2026”; *the PDF author field reads “ChatGPT; prepared for Vladimir Reshetnikov and the ProveIt project”*. The title page is dated 7 October 2026; the archive arrived at 22:07 on 6 October 2026 (Pacific time). Archive `gowers_asymptotic_constants.zip` (arrival `0d93c1eac`), manuscript `gowers_asymptotic_constants.tex` (2,633 lines, 35 pages), pin `d7b83fa1f` (its `SOURCE_AUDIT.md`, shipped as `63-asymptotic-constants-SOURCE_AUDIT.md`, records three blobs with their SHA-256 values and the predecessor manuscript of its Section 6 by SHA-256). Printed below: the abstract, the statement at the foot of its title page and Section 1; Sections 2–5 open the Part III chapter on the odd-order constants (Section III.94), Section 6 forms a Part IV chapter on Proposition 17.7 (Section IV.67), Section 7 and Appendix A are in Part V. Its macro `\T` (the circle $\mathbb{R}/\mathbb{Z}$, not this report’s $\mathbb{T}$) is renamed `\Tquot`; its ‘question’ environment (“Research question”, on the shared counter) is set as this report’s ‘Research question’; its two figures are shipped as `figures/63-asymptotic-constants-norm_upper_bounds.pdf` and `figures/63-asymptotic-constants-localization_retention.pdf`.

Abstract of source 63 (as delivered). We prove two asymptotically sharp results for intermediate problems arising from Gowers’ proof of Szemerédi’s theorem. For real centered functions on arbitrary finite abelian groups of odd order, the optimal constant in $\|f\|_{U^2}^4 \leq c_d \|f\|_{U^d}^4$ tends to exactly $1/3$, resolving the limiting-constant question posed in the inspected ProveIt reports. We give the uniform rate $c_d \leq 1/3 + 102^{-\lfloor d/8 \rfloor}$ for $d \geq 80$. The argument uses a sharp bounded inequality with the exact order- N coefficient $1/3 - 4/(3N^2) + 1/N^3$, a Fourier-tail estimate, positive smoothing, and uniform recovery of a supremum norm. It also identifies the asymptotically sharp quartic coefficient in centered cube-count upper expansions.

For local phase removal, an exact positive measure on the faces of a box improves the preceding sliding-window localization coefficient. The retained factor is $2^{k+1}/((k+2)!D_k)$, where $1 \leq D_k \leq \exp(2/[3(k+2)])$. A matching geometric ceiling proves that the best local coefficient is asymptotic to $2^{k+1}/(k+2)!$. We supply a finite theorem with explicit rounding factors, extend the cover identity to positive real powers, and formulate further extremal, stability, and formalization problems. Complete proofs, exact supplementary checks, source provenance, and internal review notes accompany the article.

Research contribution for integration into the Ramsey collection. Theorems are stated with their full scope; fixed-degree and global questions are distinguished explicitly.

Gowers’ proof of Szemerédi’s theorem combines uniformity norms, additive structure, polynomial phases, and localization to arithmetic progressions [1]. Its intermediate estimates give several well-defined extremal problems even when improvements at those steps do not change the final asymptotic bound for progression-free sets. This article addresses two such problems: the comparison between second and high-order uniformity norms for real centered functions, and the amount of local cube mass that survives removal of a multilinear derivative phase.

The first problem is uniform over all finite abelian groups of odd order. Earlier reports in the ProveIt research collection established an exact answer at order three and left the high-order limit between $1/3$ and $1/2$ [258]. We prove that the limit is $1/3$, and give an explicit convergence rate. A short sharp inequality for bounded functions supplies the endpoint; the main work is to pass from that inequality to an estimate uniform over groups of arbitrary size. The second problem concerns Proposition 17.7 of Gowers’ paper. An earlier report replaced the enumeration of carry patterns by sliding windows and positive covers [259]. We construct a more efficient cover and prove that its coefficient approaches the universal geometric ceiling as the degree increases.

F.68.1 The odd-order norm limit

For $d \geq 2$, define

$$c_d^{\text{odd}} = \sup_{\substack{|G| \text{ odd, } G \text{ finite abelian} \\ f: G \rightarrow \mathbb{R}, \mathbb{E}f=0, f \neq 0}} \frac{\|f\|_{U^2(G)}^4}{\|f\|_{U^d(G)}^4}.$$

All norms use probability normalization. Theorem III.97.1 proves

$$\boxed{\lim_{d \rightarrow \infty} c_d^{\text{odd}} = \frac{1}{3}}, \quad \frac{1}{3} \leq c_d^{\text{odd}} \leq \frac{1}{3} + 10 \cdot 2^{-\lfloor d/8 \rfloor} \quad (d \geq 80). \quad (\text{F.68.1})$$

The supremum includes both G and f . Establishing convergence for each fixed group would leave open the possibility that increasingly large groups retain a larger ratio. The Fourier-tail estimate, positive smoothing, and finite-spectrum recovery proved in Section III.97 eliminate that possibility.

The bounded-function ingredient is itself exact. If $N \geq 3$ is odd, $|G| = N$, and $f : G \rightarrow \mathbb{R}$ is centered, then

$$\|f\|_{U^2}^4 \leq \left(\frac{1}{3} - \frac{4}{3N^2} + \frac{1}{N^3} \right) \|f\|_{\infty}^4. \quad (\text{F.68.2})$$

For each N , a cyclic square wave attains equality. Thus the coefficient is the exact optimum among groups of that order. Its uniform limit is the sharp constant $1/3$. The proof is based on the autocorrelation constraint

$$1 + R_f(2h) \geq 2|R_f(h)| \quad (\|f\|_{\infty} \leq 1),$$

together with measure preservation under doubling. It also yields a compact-group theorem and an exact defect identity.

The norm result has a direct counting consequence. The quartic term of the centered d -cube expansion at background density δ is

$$P_d \delta^{2^d - 4} \|f\|_{U^2}^4, \quad P_d = \frac{6^d - 2 \cdot 4^d + 2^d}{8}.$$

Consequently the optimal quartic coefficient measured in $\|f\|_{U^d}^4$ is asymptotic to $P_d/3$. Section III.98 retains the higher-order remainder explicitly and explains the precise small-perturbation sense of this assertion.

F.68.2 Efficient phase removal

For the local problem, the natural geometric coefficient is

$$v_k = \frac{2^{k+1}}{(k+2)!}.$$

It is the leading normalized number of $(k+1)$ -cubes inside a long interval. We use the same local output normalization throughout: if a partition has M cells, each of size at most U , divide its total $(k+1)$ -cube mass by MU^{k+2} .

For an input box of odd width m in a prime cyclic group of order $N \geq m^2$, under its stated derivative-energy and width hypotheses, Theorem IV.68.5 gives a phase of degree at most $k+1$

and a partition with cell lengths $\lceil m/(3k) \rceil$ or $\lceil m/(3k) \rceil + 1$. As $m \rightarrow \infty$ for fixed k , the retained coefficient is

$$\frac{v_k}{D_k}, \quad D_k = \sum_{j=0}^k \binom{k}{j} \left(\frac{2}{3k}\right)^j \frac{(k-j+2)!}{(k+2)!}. \quad (\text{F.68.3})$$

If c_k^{loc} denotes the optimal asymptotic coefficient for this output class, then

$$e^{-2/[3(k+2)]} v_k \leq \frac{v_k}{D_k} \leq c_k^{\text{loc}} \leq v_k, \quad \boxed{c_k^{\text{loc}} \sim v_k}. \quad (\text{F.68.4})$$

The local-scale limit precedes the degree limit. The exact finite theorem records all rounding losses and its width condition, so the asymptotic statement does not conceal restrictions on finite inputs.

The geometric mechanism is useful independently of phase removal. We put explicit positive densities on all faces of a box, with codimension- j density $(k-j+2)!w^j/(k+2)!$. Convolution with $(w - \|x\|_1)_+^2$ is given by a sum of nonnegative terms, the first of which is the full integral of the kernel. Mean-preserving rounding passes the cover to lattice points without increasing its mass. The same identity extends to every positive real power of the kernel; the lattice statement holds for powers at least one.

F.68.3 Scope and relation to the source collection

The comparison baseline is the inspected ProveIt snapshot **d7b83fa1fe60**, together with the separate manuscript *Local Phase Removal and Cubic Rigidity in Gowers’ Szemerédi Framework* [259]. The source audit in Appendix V.238 records the relevant files, report numbers, and provenance. In that collection, Source 29 gave $c_d^{\text{odd}} \leq B_d$ with $B_d \downarrow 1/2$, while the lower construction approached $1/3$. The limiting-constant question posed there is answered by (F.68.1). The new cover closes the constant-factor asymptotic gap in the preceding local phase-removal coefficient.

The distinctions between these statements matter. The exact values of c_d^{odd} at fixed $d \geq 4$ remain undetermined here. The inspected later report 45 improves the fourth-order upper bound to $(832/4795)^{1/4}$; this article’s high-order result is compatible with that sharper fixed-order estimate and does not replace it. Likewise, (F.68.4) determines the first-order behavior of the optimal local coefficient but leaves its fixed-degree value and its next relative term open. Finally, we do not propagate these improvements through all the subsequent density-increment steps of Gowers’ proof. No new final quantitative bound for Szemerédi’s theorem is asserted.

The article presents complete mathematical proofs of its principal claims, with supplementary exact finite checks. The development and separate proof reviews were carried out by ChatGPT instances; the review notes are included so their scope can be assessed. They are not external peer review or a formal proof-assistant certificate. The stated research advance is relative to the inspected project questions and comparison theorems. A targeted literature search is recorded in the accompanying audit, without a claim of historical priority over all published work.

[write] 7 October 2026; sources 60–69. *Identifications and credit within this report.* [258] is this repository at **d7b83fa1f**, an ancestor of HEAD; the transcription (**41085a0**), this article (**32087ba**, the sixth write) and the README (**689dbd3**) recorded in its source audit, shipped as **63-asymptotic-constants-SOURCE_AUDIT.md**, were verified by blob and SHA-256. Its “Source 29” and “report 45” are this report’s (Sections III.66 and III.89); the fourth-order bound it quotes from source 45 has since been improved by source 49 to $(4/29)^{1/4}$. [259], the predecessor credited for its Section 6, was *never delivered* to this repository (Section F.4). *Not cited*: source 16, whose sliding-window localization and local ceiling $2^{k+1}/(k+2)!$ its Section 6 uses (Section IV.67); and source 60, which arrived with it and proves a weaker form of its limit theorem. Its three verifiers were rerun on copies at the write.

F.69 Source 64. The problem, the results, and their position in the literature

Source 64, delivered Section 1.

Added in the eleventh write, 7 October 2026. Delivered title block. “Binomial separation and endpoint bounds for uniform progression counts / A six-variable digit construction, an exact nine-copy transfer, and obstructions to perfect cyclic codes / Research manuscript prepared for the ProveIt project / 6 October 2026”; *the PDF author field reads the same; neither the manuscript nor its delivery files make a statement on AI assistance.* Archive `gowers_binomial_separation_research.zip` (arrival `0d93c1eac`), manuscript `article.tex` (1,268 lines, 19 pages); *no pin:* its `SOURCE_AUDIT.md`, shipped as `64-binomial-separation-SOURCE_AUDIT.md`, records the blob `987daf1` of `Section04.lean` and two “search snapshots” `17123532a` and `176e155e5`, “not a claim that the entire repository was checked at either immutable commit”. Printed below: the abstract, its status statement and Section 1; Sections 2–8 form a chapter of Part III on Gowers’s Section 4 conjectures (Section III.100), Sections 9–11 are in Part V. It uses `cleveref` (options `capitalize`, `nameinlink`, `noabbrev`); its references are written out with the name of the statement referred to, capitalized. *Its delivered PDF names every statement on the shared counter “Theorem”* (for example “Theorem 2.4” for Proposition 2.4 and “Theorem 4.1” for Lemma 4.1), as source 57’s did; the names printed here are those of the statements. Its one-argument macro `\repo` (a path in typewriter type, not this report’s ProveIt) is renamed `\repopath`; its ‘question’ environment (“Question”, on the shared counter) is set in this report’s ‘Question’ style.

Abstract of source 64 (as delivered). We study the quantitative separation between uniformity and arithmetic-progression counts that underlies the conjectures in Section 4 of Gowers’s 2001 paper. The alphabet $D = \{0, 1, 4, 13\}$ makes the map $(a, b, c) \mapsto a + 5b + 10c$ injective modulo 80. A digit-lifting argument therefore constructs sets with no nontrivial solutions to the six-binomial equation at reciprocal growth exponent $b_6 = \log_4 80 = 3.1609640474\dots$, improving the exponent 5 supplied by the general construction recorded by Deng–Tidor–Zhao. We give an exact transfer from a symmetric-progression-free r -coloring of an interval of length N to a $9r$ -coloring of the circle: the symmetric k -progression probability is exactly $1/(9(k-1)N)$, independently of the number of colors. Combined with binomial-separating labels, this yields an explicit torus construction with exactly computable progression density. Using the prime-indexed colorings of Shi–Dong, rather than fixing and tensoring one coloring, gives the endpoint bounds $\rho_4(\alpha) \ll \alpha^5$ and $\rho_6(\alpha) \ll \alpha^{5+9/\log_4 80}$, without an epsilon loss in the exponents. We also prove that no non-singleton cyclic digit alphabet for the weights $(1, 5, 10)$ can attain the counting limit $|D|^3 = B$, and give an exact Chevalley–Warning dimension barrier for degree- $(k-1)$ polynomial colorings. All claimed improvements have mathematical proofs; finite checks are supplied separately. The conjectural value $\beta_6 = 3$ and the superpolynomial-decay questions remain open here.

Status and scope. This is a research manuscript, not a declaration of formal verification or of publication priority. The improvements established below are compared with specific inspected sources. The Shi–Dong coloring theorem is credited as a 2026 preprint input. The general analytic realization theorem is credited to Deng–Tidor–Zhao; an independent prime-modulus proof sufficient for the present indicator constructions is included. No improvement to the general quantitative Szemerédi theorem or to classical graph Ramsey numbers is claimed.

F.69.1 Uniformity one order below the counting threshold

For $G = \mathbb{Z}/M\mathbb{Z}$ and $f : G \rightarrow \mathbb{C}$, put

$$\Lambda_k(f) = \mathbb{E}_{x,d \in G} \prod_{i=0}^{k-1} f(x + id). \quad (\text{F.69.1})$$

All progressions are ordered, and the term $d = 0$ is included. Its contribution is negligible as $M \rightarrow \infty$ for the fixed-density asymptotic questions here. We use normalized Gowers norms:

$$\|f\|_{U^s(G)}^{2^s} = \mathbb{E}_{x, h_1, \dots, h_s \in G} \prod_{\omega \in \{0,1\}^s} \mathcal{C}^{|\omega|} f(x + \omega \cdot h), \quad (\text{F.69.2})$$

where $\mathcal{C}$ denotes complex conjugation.

We use the definition of ρ_k in [37, Definition 1.11]: $\rho_k(\alpha)$ is the largest u such that, for every $\varepsilon > 0$, there are $\eta > 0$ and M_0 for which

$$M \geq M_0, \quad A \subseteq \mathbb{Z}/M\mathbb{Z}, \quad \|\mathbf{1}_A - \alpha\|_{U^{k-2}} \leq \eta \implies \Lambda_k(\mathbf{1}_A) \geq u - \varepsilon. \quad (\text{F.69.3})$$

In particular, a sequence of examples along primes suffices for an upper bound on this quantity, even though the definition quantifies over all sufficiently large moduli.

Gowers's Section 4 conjectures concerned lower bounds of random-set order under one less degree of uniformity than the standard counting theorem requires [1]. Gowers's later construction already gave a counterexample at length four [260]. The 2026 Shi–Dong preprint gives the corresponding negative conclusion for every even length [33]. Those disproofs are *not* contributions claimed here. Our target is the sharper quantitative question: how small can the progression count be, and which parts of the coloring-to-uniform-set construction lose exponent?

F.69.2 Main quantitative statements

Set

$$b_6 = \frac{\log 80}{\log 4} = 3.160964047443681 \dots, \quad \vartheta_6 = 5 + \frac{9}{b_6} = 7.847232636916081 \dots \quad (\text{F.69.4})$$

Here b_6 is the exponent furnished by our construction; β_6 below continues to denote an extremal parameter, not a quantity asserted to equal b_6 .

Theorem F.69.1 (Principal consequences). *The following statements hold.*

- (i) *For every positive integer r and every integer $m > 256r^{b_6}$ there is an r -element set $S \subseteq \mathbb{Z}/m\mathbb{Z}$ such that*

$$x_0 - 5x_1 + 10x_2 - 10x_3 + 5x_4 - x_5 = 0, \quad x_i \in S, \quad (\text{F.69.5})$$

forces $x_0 = x_5$, $x_1 = x_4$, and $x_2 = x_3$. Consequently $3 \leq \beta_6 \leq b_6$, where β_6 is the reciprocal asymptotic growth exponent for sets avoiding nontrivial six-binomial solutions.

- (ii) *There are constants $C_4, C_6, \alpha_0 > 0$, independent of α , such that*

$$\rho_4(\alpha) \leq C_4 \alpha^5, \quad \rho_6(\alpha) \leq C_6 \alpha^{\vartheta_6} \quad (0 < \alpha \leq \alpha_0). \quad (\text{F.69.6})$$

For each fixed such α , these bounds are realized by deterministic indicator sets along primes, with density tending to α and centered U^{k-2} norm tending to zero.

- (iii) *If $D \subseteq \mathbb{Z}/B\mathbb{Z}$, $|D| = r > 1$, and $(a, b, c) \mapsto a + 5b + 10c$ is injective on D^3 , then $B > r^3$. Thus no one fixed cyclic digit alphabet attains reciprocal exponent 3.*

Part (i) is proved in Section III.101; part (ii) follows from the exact transfer in Sections III.103 and III.104 and the scale-selection argument in Section III.105; part (iii) is proved in Section III.102. An additional barrier in Section III.107 shows that the dimension $k^2/4$ in the scalar degree- $(k-1)$ polynomial stage of the Shi–Dong construction cannot be increased within that precisely specified model.

F.69.3 What is strengthened, and what is inherited

The comparison points are deliberately specific. The six-binomial construction improves the universal exponent $k - 1 = 5$ in [37, Lemma 3.3 and Definition 3.4], and is a six-variable counterpart of that paper’s special four-variable label lemma. The circle construction retains only full mirror symmetry, which is sufficient when the labels force mirror pairings. This removes the need to eliminate additional monochromatic patterns that enter the more general reduction. The endpoint argument uses all scales of the available prime family, so constants depending on a single chosen prime never turn into an exponent loss.

For four-term progressions the resulting exponent 5 is the endpoint of the $5 - \varepsilon$ family stated in [33, Corollary 1.4]. For six terms the combination with the generic exponent 5 yields a limiting exponent $34/5$ before an epsilon loss; replacing the labels and avoiding that loss gives $\vartheta_6 > 34/5$. Since $\alpha < 1$, the larger exponent is the stronger upper bound. These are not statements that the displayed exponents are optimal.

The language of injective weighted sums is the standard language of *Sidon sets for linear forms* [261]. Digital constructions are not new as a general method; see also the distinct setting of greedy weighted Sidon sets in [262]. The proposed contributions are the explicit six-binomial parameters, the exact transfer and its endpoint use, and the precise obstructions proved here. The literature and repository searches recorded in the accompanying source audit are not an exhaustive priority

[write] **7 October 2026; sources 60–69.** *Identifications within this report. This source records no pin:* its source audit, shipped as `64-binomial-separation-SOURCE_AUDIT.md`, names the blob `987daf1` of `Section04.lean`, which is the file at HEAD `fb9602e55` (`conjecture_4_1`, `conjecture_4_2` at lines 30 and 39), and two search snapshots `17123532a` and `176e155e5`, both ancestors of HEAD. It is the report’s first source on Gowers’s Section 4 (Section III.100). *Not cited:* sources 10 and 14, whose exact dimension threshold contains this source’s barrier of Section III.107 (note after Theorem III.107.2); both were in this report at its search snapshots. [33] and [37] are the report’s entries for the preprint and paper it uses. Its verifier was rerun on a copy at the write.

search.

F.70 Source 65. The extremal problem and the main results

Source 65, delivered Section 1.

Added in the eleventh write, 7 October 2026. Delivered title block. “Sharp Boolean Phase Integration in Every Degree / Defect-radical slicing, exact nonintegrability thresholds, and extremal tensor rigidity in Gowers’s framework / Research manuscript prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt project / 6 October 2026”; *the PDF author field reads “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov”*. Archive `Gowers_All_Degree_Boolean_Integration.zip` (arrival `0281fcb05`, 7 October 2026, 08:30), manuscript `all_degree_boolean_integration.tex` (1,459 lines, 21 pages), pin `8c40adc24` (its `SOURCE_AUDIT.md`, shipped as `65-all-degree-boolean-SOURCE_AUDIT.md`, records the blobs of `Sections17_18.lean` and `Definitions.lean`). Printed below: the abstract, its status statement and Section 1; Sections 2–8 open the Part IV chapter on Boolean phase integration in every degree (Section IV.51), Sections 9–12 and Appendices A–B are in Part V. Its macros `\F` ($\mathbb{F}_2$, not this report’s $\mathbb{F}$), `\TT` (the circle $\mathbb{R}/\mathbb{Z}$, not this report’s $\mathbb{T}$) and `\pin` (its own pin) are set as this report’s `\Ftwo`, the same definition, as `\Tquot` and as `\pinsixtyfive`; its ‘question’ environment (“Research question”, numbered 1–9 on its own counter) is set in this report’s ‘Research question’ style, numbered within its section, which keeps the indices.

Abstract of source 65 (as delivered). For a symmetric d -linear frequency tensor T on $\mathbb{F}_2^n$, let $\mathcal{E}_T(f)$ be its normalized selected derivative Fourier energy. We prove that the largest

possible energy of a nonintegrable symmetric tensor is exactly $1 - (d + 1)/2^d$, for every $d \geq 3$ and $n \geq 2$. This closes the septic and all-degree questions in the preceding project manuscript, which established the universal value only through degree six. The key is an exact identification: the directions producing integrable contractions are precisely the radical of the repeated-variable defect. A noncanonical defect has radical codimension at least two, so at least three quarters of its contractions remain nonintegrable. Slicing and induction then give the strictly smaller cap $1 - 3d/2^{d+1}$. Consequently every extremal tensor, in every degree, is a canonical two-coordinate obstruction modulo an integrable tensor. We give an explicit recovery algorithm, count these obstruction classes, classify all their integrable restrictions, and obtain a dimension-independent rigidity window and robust phase-removal criteria. Full proofs include the needed cubic spectral bound and the prior canonical energy formula. The integration criterion and canonical-family result are credited as existing ingredients; the all-degree closure and its rigidity consequences are the contributions of this manuscript.

Status and scope. Unrefereed research manuscript with exact finite regression checks, not a proof-assistant verification. The resolved questions are explicitly identified project questions. Literature priority is not asserted. This is a sharp characteristic-two local phase-integration result, not a new end-to-end bound for Szemerédi’s theorem.

F.70.1 Normalization and integrability

Let $V = \mathbb{F}_2^n$. All averages are uniform and normalized. For a function $f : V \rightarrow \mathbb{C}$ and a direction $h \in V$, write

$$\partial_h f(x) = f(x + h) \overline{f(x)}, \quad \Delta_h P(x) = P(x + h) - P(x), \quad e(t) = \exp(2\pi i t).$$

A tensor in this article is genuinely multilinear in all its slots, not merely a polynomial that is affine in each slot. Write $\text{SymMult}^d(V)$ for symmetric multilinear maps $V^d \rightarrow \mathbb{F}_2$. For a d -linear tensor T define

$$\mathcal{E}_T(f) = \mathbb{E}_{h_1, \dots, h_{d-1}} \left| \mathbb{E}_x \partial_{h_1} \cdots \partial_{h_{d-1}} f(x) (-1)^{T(h_1, \dots, h_{d-1}, x)} \right|^2, \quad \mathcal{M}(T) = \max_{\|f\|_\infty \leq 1} \mathcal{E}_T(f). \quad (\text{F.70.1})$$

The last slot evaluates the selected frequency. The maximum exists by compactness of a finite product of closed complex unit discs. In this normalization $0 \leq \mathcal{E}_T(f) \leq 1$ for $\|f\|_\infty \leq 1$.

A map $P : V \rightarrow \mathbb{R}/\mathbb{Z}$ is a nonclassical polynomial of degree at most d when all its $(d + 1)$ -fold additive differences vanish. A symmetric tensor T is *integrable* when some such P satisfies

$$\Delta_{h_1} \cdots \Delta_{h_d} P(x) = \tfrac{1}{2} T(h_1, \dots, h_d) \pmod{1} \quad \text{for all } x, h_1, \dots, h_d. \quad (\text{F.70.2})$$

The integrable tensors form a linear subspace $\mathcal{I}_d(V)$. Circle-valued phases are essential here: restricting the primitive to an $\mathbb{F}_2$ -valued classical polynomial would change the problem.

For independent $u, v \in V^*$ put

$$C_d^{u,v}(h_1, \dots, h_d) = \sum_{i=1}^d v(h_i) \prod_{j \neq i} u(h_j), \quad b_d = 1 - \frac{d+1}{2^d}. \quad (\text{F.70.3})$$

Call a tensor *canonical modulo integration* when it equals $C_d^{u,v} + S$ with u, v independent and $S \in \mathcal{I}_d(V)$.

F.70.2 The universal theorem

Theorem F.70.1 (Sharp all-degree symmetric threshold). *For every $d \geq 3$ and $n \geq 2$,*

$$\boxed{\max_{T \in \text{SymMult}^d(V) \setminus \mathcal{I}_d(V)} \mathcal{M}(T) = 1 - \frac{d+1}{2^d}.} \quad (\text{F.70.4})$$

Moreover, for a nonintegrable symmetric tensor T ,

$$\mathcal{M}(T) = b_d \iff T = C_d^{u,v} + S \quad \text{for independent } u, v \in V^*, S \in \mathcal{I}_d(V). \quad (\text{F.70.5})$$

If $\mathcal{E}_T(f) > b_d$, then T is integrable. For any primitive P_T ,

$$\mathcal{E}_T(f) = \|fe(-P_T)\|_{U^d}^{2^d}. \quad (\text{F.70.6})$$

The strict inequality cannot be replaced by a non-strict one.

[write] **7 October 2026; sources 60–69.** *Proved twice* (Section IV.51): source 69's Theorem F.74.1 is the same statement for the maximum, with the equality classification for $d \geq 4$. *Checked at the write* (exact rationals): the table above, b_d for $3 \leq d \leq 10$, the noncanonical caps $1 - 3d/2^{d+1}$ for $4 \leq d \leq 10$ and $1/4$ at $d = 3$, the gaps $(d-2)/2^{d+1}$, source 47's caps $1 - \min\{d+1, 7\}/2^d$ and the gain $(d-6)/2^d$.

Here the Gowers quantity is normalized by $\|g\|_{U^d}^{2^d} = \mathcal{E}_0(g)$. The conclusion concerns the prescribed tensor T , not the construction of a tensor from an arbitrary large Gowers norm. Full symmetry includes the final frequency-evaluation slot.

F.70.3 The quantitative separation behind the theorem

For $d \geq 3$, the repeated-variable defect is

$$\Phi_T(z_1, \dots, z_{d-3})(a, b) = T(a, a, b, z_1, \dots, z_{d-3}) + T(a, b, b, z_1, \dots, z_{d-3}). \quad (\text{F.70.7})$$

For $d = 3$ denote this alternating bilinear form by Ω_T . For $d \geq 4$ define

$$K_T = \{z \in V : \Phi_T(z, z_2, \dots, z_{d-3}) = 0 \text{ for every } z_2, \dots, z_{d-3}\}, \quad r(T) = \text{codim } K_T. \quad (\text{F.70.8})$$

When $d = 4$ there are no remaining frozen variables. The zero on the right means the zero alternating form in a, b .

Theorem F.70.2 (Radical-sensitive cap and strict noncanonical gap). *Let $d \geq 4$ and $T \in \text{SymMult}^d(V)$ be nonintegrable. If $r(T) = 1$, then T is canonical modulo integration and $\mathcal{M}(T) = b_d$. If $r(T) \geq 2$, then*

$$\boxed{\mathcal{M}(T) \leq 1 - (1 - 2^{-r(T)}) \frac{d}{2^{d-1}} \leq 1 - \frac{3d}{2^{d+1}} = b_d - \frac{d-2}{2^{d+1}}.} \quad (\text{F.70.9})$$

For $d = 3$, the exact formula is $\mathcal{M}(T) = 2^{-\text{rank } \Omega_T/2}$; the noncanonical cap is $1/4$.

The bound in (F.70.9) is *not* asserted to be the optimal second extremal value. Its strict separation from b_d is enough to classify all endpoint tensors.

F.70.4 The project question that is closed

The manuscript *Exact Boolean Obstruction Energies* [264] proved the canonical-family formula $\mathcal{M}(C_d) = b_d$ in every degree, but its universal argument gave only

$$b_d \leq \max_{T \notin \mathcal{I}_d(V)} \mathcal{M}(T) \leq 1 - \frac{\min\{d+1, 7\}}{2^d}.$$

Its Research Questions 1 and 2 asked for the septic value and the universal all-degree value. Its Research Question 4 also asked for the quintic and sextic endpoint tensors and functions. Theorem F.70.1 answers the first two questions and the *tensor* part of the fourth, in fact in every degree. Classification of all maximizing functions remains a different question.

Degree	Prior universal cap	Sharp value here	Noncanonical cap here	Gap below sharp value
3	1/2	1/2	1/4	1/4
4	11/16	11/16	5/8	1/16
5	13/16	13/16	49/64	3/64
6	57/64	57/64	55/64	1/32
7	121/128	15/16	235/256	5/256
8	249/256	247/256	61/64	3/256
9	505/512	251/256	997/1024	7/1024
10	1017/1024	1013/1024	1009/1024	1/256

The all-degree gain over the preceding universal cap is $(d-6)/2^d$ for $d \geq 7$. More significantly, there is now a complete extremal theorem rather than a sequence of isolated low-degree support estimates.

F.70.5 What is an existing ingredient and what changes

The Boolean nonclassical integration criterion is established background [150, Definition 3.1 and Proposition 3.5]. The cubic spectral law, the shear argument, and the canonical-family formula were developed in the preceding project drafts [263, 264]. We reprove exactly the pieces needed, so the new conclusions do not depend on accepting an unchecked earlier proof.

The new step is to count *integrable contractions*, not merely the support of Φ_T . Their set is exactly K_T . Radical codimension one is the already understood canonical case; radical codimension at least two leaves at least three quarters of the slices nonintegrable. The sharp theorem in degree $d-1$ therefore supplies a stronger estimate than a degree-independent support bound. This also gives the endpoint classification, its finite rigidity window, and the reconstruction and counting consequences below.

[write] 7 October 2026; sources 60–69. *Identifications within this report.* [264] (“the preceding project manuscript”) is source 47 (Section IV.42; its comparison pin 5c9a442d2), [263] is source 43 (Section IV.25), and [265] is this repository at 8c40adc24, an ancestor of HEAD; Sections17_18.lean and Definitions.lean have the recorded blobs at HEAD fb9602e55. Source 47’s Research questions V.176.1, V.176.2 and V.176.4 are the ones this source names. *Not cited, and not citing it:* source 69, which proves the main theorem independently on the same day (Section IV.51). The verifier was rerun on a copy at the write.

F.71 Source 66. Scope, previous results, and principal contributions

Source 66, delivered Section 1.

Added in the eleventh write, 7 October 2026. Delivered title block. “Partition Normal Forms for Small Schur Defect / Sharp endpoint bounds and the first four additive-energy levels / A local strengthening of the ordered-energy step in Gowers’s proof / Research manuscript prepared with ChatGPT for Vladimir Reshetnikov / 6 October 2026”; *the PDF author field reads the same*. Archive `Gowers_Endpoint_Partitions.zip` (arrival 0281fcb05), manuscript `article.tex` (1,637 lines, 24 pages); *no pin*: its `SOURCES.md`, shipped as `66-endpoint-partitions-SOURCES.md`, records three blobs as “audit anchors” (the transcription, this directory’s README and source 45’s energy review) and says that “no uniform snapshot of the whole repository is claimed”. Printed below: the abstract, its keywords and Section 1; Sections 2–11 stand in the Part II chapter on the energy levels after source 68 (Section II.296), Sections 12–14 and Appendices A–C are in Part V. Its macros `\E` (roman E, the additive energy, not this report’s expectation $\mathbb{E}$), `\D` (roman D, the energy defect, not this report’s Δ), `\ep` (roman e, the endpoint defect), `\card` ($|\cdot|$, not this report’s operator) and `\wt` (“wt”, not this report’s “wid”) are set as `\Erm`, `\Drm`, this report’s `\e` (the same definition), `\cardbars` and `\wtop`; its ‘question’ and ‘conjecture’ environments (on the shared counter) are set as this report’s ‘Question’ and a ‘Conjecture’ environment added here; its delivered `data/partition_spectrum.csv` has CRLF line ends and is shipped byte for byte.

Abstract of source 66 (as delivered). For a finite real set A of size m , write $E(A) = |\{(a, b, c, d) \in A^4 : a + b = c + d\}|$ and $M_m = (2m^3 + m)/3$. Gowers’s Section 7 recalls the sharp inequality $E(A) \leq M_m$, with arithmetic progressions as the equality cases. We refine this endpoint through a classification of small Schur defects. If $0 < c_1 < \dots < c_n$ and at most r of the differences $c_i - c_j$ with $j < i$ are absent from the set, then, for $n \geq 3r + 1$, there is a $d > 0$ such that $c_i = d(i + \lambda_i)$, where the λ_i are nonnegative, nondecreasing integers and their sum is exactly the defect. Consequently, defect r has precisely $p(r)$ positive-dilation classes. Reflection gives partition-indexed, bounded-hole normal forms for endpoint extensions of ordered sets. Their full energy is an explicit affine function of m , and the sharp conditional bound is $E(A) \leq M_m - 4r(m - 1 - r)$ when the last endpoint has defect r and $m - 1 \geq 3r + 1$. Equality forces a single hole at position r . Using the cases $r \leq 3$, we classify all third-level extremizers for $m \geq 8$ and all fourth-level extremizers for $m \geq 11$. In particular, the first four energy levels for $m \geq 11$ consist of eight oriented arithmetic templates. Every set in these levels lies in an $(m + 2)$ -term progression and contains an $(m - 2)$ -term progression. The energy classifications extend to torsion-free abelian groups. All general statements have elementary written proofs; exact finite checks accompany the article. Publication priority is not established, and no improved global bound for Szemerédi’s theorem is claimed.

Keywords. Additive energy; Schur triples; inverse theorems; arithmetic progressions; integer partitions; Gowers’s proof.

F.71.1 The precise interface with Gowers’s argument

Immediately before Theorem 7.2 of [1], Gowers considers the invariant

$$E(A) = \sum_t |A \cap (A + t)|^2 = \left| \left\{ (a, b, c, d) \in A^4 : a + b = c + d \right\} \right|. \quad (\text{F.71.1})$$

For an m -element subset of the integers, the maximum is

$$M_m = m^2 + 2 \sum_{j=1}^{m-1} j^2 = \frac{2m^3 + m}{3}, \quad (\text{F.71.2})$$

and equality characterizes arithmetic progressions. The ensuing Balog–Szemerédi statement and Proposition 7.3 extract a large structured subset from a set with energy at least $c_0 m^3$. The graph furnished by the preceding inverse step is one reason this energy invariant is useful.

Our subject is the fine structure very near the *ordered* maximum (F.71.2). We retain exact integer deficits instead of replacing them by density parameters. This is not the finite-subgroup endpoint m^3 : a torsion-free set has a different extremal geometry. The results below therefore apply directly to ordered sets and torsion-free lifts, not automatically to an arbitrary subset of a finite cyclic group.

F.71.2 What is prior, and what is added here

The repository’s source-45 energy proof review [266] records the ordered-defect identity, an elementary proof of the first energy gap, the complete second-level classification, a unit-Schur-defect lemma, and the sharp third-level exclusion bound. Those facts are credited and reproved here so that the new arguments have no dependency on an unverified repository theorem. The first gap is also implicit in the published equality theory for Pollard’s inequality [190]; it is not presented as a new inequality.

The additions developed in this article are the all- r partition normal form, its exact boundary-hole energy law and sharp conditional envelope, and the complete third- and fourth-level equality analyses. The latter are refinements of the third-level bound recorded in [266], not merely new examples attaining that bound. We also give a full-set inverse conclusion and an affine graph consequence in the corresponding energy window.

A bounded literature search and a reading of the relevant repository materials do not establish publication priority. In particular, the absence of a matching statement in the sources inspected is not evidence that no equivalent theorem exists elsewhere. The mathematical claims are the explicitly stated theorems and their proofs; claims of priority are deliberately more limited. The manuscript is AI-assisted, unrefereed, and not proof-assistant verified. The included program checks finite instances only.

F.71.3 Main statements in one place

Throughout, $[u, v]_{\mathbb{Z}} = \{u, u + 1, \dots, v\}$ for integers $u \leq v$. An oriented affine copy means $b + dS$ with $d > 0$. Keeping the orientation is necessary when discussing the *last* endpoint.

For a positive ordered set $C = \{c_1 < \dots < c_n\}$, define

$$\sigma(C) = \#\{(i, j) : 1 \leq j < i \leq n, c_i - c_j \notin C\}. \quad (\text{F.71.3})$$

Theorem II.302.1 proves that, when $n \geq 3r + 1$ and $\sigma(C) \leq r$,

$$c_i = d(i + \lambda_i), \quad 0 \leq \lambda_1 \leq \dots \leq \lambda_n, \quad \lambda_i \in \mathbb{Z}, \quad \sum_i \lambda_i = \sigma(C). \quad (\text{F.71.4})$$

For exact defect r , the positive entries form an integer partition of r . Thus there are exactly $p(r)$ possibilities up to positive dilation, independently of n in this range.

For $A = B \cup \{a\}$ with $a > \max B$, the *endpoint defect* is

$$e(B, a) = \#\{(z, x) \in B^2 : z < x, a - x + z \notin B\}. \quad (\text{F.71.5})$$

This is not the total energy defect. Theorem II.305.2 proves that, if $|A| = m$, $e(B, a) = r \geq 1$, and $m - 1 \geq 3r + 1$, then

$$E(A) \leq M_m - 4r(m - 1 - r). \quad (\text{F.71.6})$$

Equality holds precisely for an oriented affine copy of $[0, m]_{\mathbb{Z}} \setminus \{r\}$. A stronger statement classifies every possible A under the endpoint hypothesis, not only equality.

Finally, Theorem II.308.2 gives the four largest distinct values of $E(A)$, for $m \geq 11$:

$$M_m, \quad M_m - 4(m - 2), \quad M_m - 8(m - 3), \quad M_m - 8m + 20. \quad (\text{F.71.7})$$

Their equality cases are respectively one, two, three, and two oriented arithmetic templates. In particular,

$$E(A) \geq M_m - 8m + 20 \implies A \text{ lies in an } (m + 2)\text{-term arithmetic progression.} \quad (\text{F.71.8})$$

The same hypothesis produces an $(m - 2)$ -term progression inside A . It is a much stronger conclusion than a positive-fraction extraction, but it requires a correspondingly much stronger hypothesis.

[write] **7 October 2026; sources 60–69.** *Identifications within this report.* [266] is source 45’s energy review, shipped as `45-rigidity-gaps-review_energy.md` (blob 7984f49, unchanged at HEAD fb9602e55); the README it read is the seventh write’s (1ec6351), the transcription blob 41085a0. *This source records no pin* (Section F.4). *Not cited, and not citing it:* source 57, whose partition normal form and third energy level it reaches by another route the same day (notes after Theorems II.302.1 and II.307.1); source 57 was not in the repository when this source was written. [190] is the report’s entry for the inverse-Pollard paper (also cited by source 45). Its verifier was rerun on a copy at the write.

F.72 Source 67. Research position and the exact improvements

Source 67, delivered Section 1.

Added in the eleventh write, 7 October 2026. Delivered title block. “Exact Local Profiles and Structural Barriers in Gowers’ Szemerédi Framework / Unequal Fourier budgets, characteristic-three rigidity, optimal interpolation lists, and norm layers without a degree cap / Research manuscript prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt Ramsey project / 7 October 2026”; *the PDF author field reads “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov”*. Archive `Gowers_Exact_Local_Profiles_TeX_and_PDF.zip` (arrival 0281fcb05), manuscript `gowers_exact_local_profiles.tex` (3,295 lines, 44 pages), the single-file assembly of its front matter, four section files and end matter by its `code/assemble.py` (shipped as `code/67-exact-profiles-assemble.py`; the assembly was reproduced byte for byte at the write); *no pin*: its `SOURCE_LEDGER.json`, shipped as `data/67-exact-profiles-SOURCE_LEDGER.json`, identifies its inputs by blob and SHA-256 and says that moving-branch paths “are not asserted to be immutable commit snapshots”. Printed below: the abstract, its keywords and Section 1; Section 2 (unequal Fourier budgets) stands in Part I after source 37’s chapter (Section I.75), Sections 3–4 form a Part II chapter on affine rigidity in characteristic three (Section II.293), Section 5 a Part II chapter at Lemma 16.10 (Section II.338), Section 6 a Part III chapter on norm layers (Section III.108), Sections 7–8 are in Part V. It defines no macros; its ‘question’ environment (“Question”, on the shared counter) is set in this report’s ‘Question’ style; its two figures are shipped as `figures/67-exact-profiles-fourier_curvature.pdf` and `figures/67-exact-profiles-rigidity_profile.pdf`. The comments of its assembly inside sections are kept; those marking the joins between section files are dropped with the section breaks.

Abstract of source 67 (as delivered). We develop four sets of exact refinements of local problems arising in Gowers’ proof of Szemerédi’s theorem and in the existing ProveIt research series. First, a supported moment problem gives an exact matrix formula for missing-arc mass

under unequal Fourier budgets, on an explicitly certified region containing neighbourhoods of every interior point of the known equal-budget branch. We determine the entire optimality polytope of the old certificate, the rank of the value Hessian, and a quadratic aligned-grid error bound. Second, for maps from any finite abelian group to an elementary abelian 3-group, the cubic affine-error profile holds throughout its positive range, with explicit color and target-rank remainders and a complete equality classification. The exact universal strict thresholds for unique affine recovery are $1/4$ for arbitrary domains and $4/9$ for elementary abelian 3-domains. The exact non-affine energy maximum in this target class has the optimal domain-size cutoff $N \geq 13$. Third, the interpolation cost at Gowers' Lemma 16.10 is reduced to the elementary symmetric polynomial of the anchor list sizes, and matching dense examples prove this cost optimal for the stated local interface, even with controlled deletion. Fourth, sums of independent field-norm blocks are classified without any degree cap in characteristic $p > 2m$: for progression lengths $2m = 4, 6, 8$, or for general length $2m$ over sufficiently large fields, their exact maximum source dimension is rm^2 for r output coordinates. The small-field eight-term case is closed by eleven explicit irreducibility certificates. Proofs, equality cases, quantitative bounds, reproducible checks, source comparisons, and a research programme are included. These are local theorems and structural obstructions; no improved global Szemerédi bound or solution of the unrestricted quintic-colouring problem is claimed.

Keywords. Additive energy; Fourier moments; affine recovery; Gowers norms; polynomial interpolation; finite-field norms; symmetrically coloured arithmetic progressions.

Gowers' proof turns information about arithmetic progressions into successive statements about Fourier coefficients, derivatives, approximately additive graphs, polynomial behaviour on structured sets, and density increments [1]. Each passage has its own extremal problem. A local constant can be sharpened only after one identifies which information is retained by that passage and what the next passage actually needs. Conversely, a sharp local obstruction can identify information that an improved argument must recover.

The ProveIt Ramsey collection already contains a substantial sequence of refinements of this framework [267]. We take that work as a baseline. In particular, the common-budget Fourier profile, the cubic $4t - 10t^2 + 6t^3$, the shared-anchor sampling improvement, and the low-degree layered norm construction are established inputs, not discoveries of this article. The contribution is to solve several remaining extremal problems attached to those inputs, and to push others to precise finite obstructions.

The word “exact” is used with a specified scope throughout. An exact Fourier formula is asserted on a stated matrix-certificate region. An exact interpolation cost is relative to explicitly given slice-list and fibre hypotheses. An exact dimension classification is for sums of independent norm blocks. A universal recovery threshold refers to the particular four-point graph-energy test and the named class of groups. These scopes are part of the theorems.

F.72.1 Four results and their relation to the baseline

A. Unequal Fourier budgets and their curvature. For probability measures on the circle, fix an open arc $I_a = \{\theta : |\theta|_o < a\}$ and bounds $|\hat{\mu}(j)| \leq \varepsilon_j$, $1 \leq j \leq R$. Let $V_{R,a}(\varepsilon)$ be the least possible mass of I_a . Source 37 obtains the exact common-budget answer in the window

$$\frac{\pi}{R+1} < a < \frac{2\pi}{R+1}, \quad a < \pi, \quad V_{R,a}(e, \dots, e) = \frac{\rho - e}{1 + \rho},$$

where

$$\rho = (R - \cot(a/2) \tan((R+1)a/2))^{-1} \quad \text{and} \quad 0 \leq e \leq \rho \quad [268].$$

The equal-budget certificate alone gives only a linear lower bound after the budgets separate. We prove that the actual optimum is

$$V_{R,a}(\boldsymbol{\varepsilon}) = -\frac{1}{(1 - \cos a) \mathbf{e}^T L(\boldsymbol{\varepsilon})^{-1} \mathbf{e}}$$

under finitely checkable strict hypotheses, where L is an R -by- R localizing matrix and $\mathbf{e}$ is the all-ones vector (Theorem I.76.2). The hypotheses hold on an open neighbourhood of every $0 < e < \rho$ on the equal-budget line.

The old certificate remains optimal on an explicitly described polytope of local dimension $\lfloor R/2 \rfloor + 1$. At every point of the strict matrix region, the Hessian of the actual value has rank $\lfloor (R-1)/2 \rfloor$. Thus for $R \geq 3$ unequal perturbations usually require a moving certificate. The rational three-frequency example in Corollary I.76.7 exhibits the exact positive quadratic gain. Theorem I.76.8 also improves the elementary grid approximation to order N^{-2} when the endpoints are aligned; it separately retains the order- N^{-1} endpoint drift.

B. The full positive cubic range in characteristic three. For $f : G \rightarrow W$, where G is finite abelian of size N and W is an elementary abelian 3-group, write

$$\mathcal{E}(f) = \#\{x_1 + x_2 = x_3 + x_4 : f(x_1) + f(x_2) = f(x_3) + f(x_4)\}, \quad \varepsilon = 1 - \mathcal{E}(f)/N^3.$$

For any affine reference ℓ , let $t = N^{-1}|\{x : f(x) \neq \ell(x)\}|$ and $\Psi(t) = 4t - 10t^2 + 6t^3$. Theorem II.294.1 proves

$$\varepsilon \geq \Psi(t) \quad (0 \leq t \leq 2/3)$$

with a positive remainder measuring multiple error colors and target rank. For $0 < t < 2/3$, equality holds exactly when $f - \ell$ is one nonzero constant on a subgroup coset and zero elsewhere. The endpoint $t = 2/3$ has a different, fully described affine equality family. The common cubic and the one-color coset examples occur in the previous reports; the enlarged range and the new projection-sensitive remainder are proved here [270, 271].

This local theorem combines with a Fourier argument to give unique nearest-affine recovery whenever $\varepsilon < 1/4$ on arbitrary finite abelian domains, or whenever $\varepsilon < 4/9$ on elementary abelian 3-domains (Theorem II.295.3). Both thresholds are best possible as universal strict thresholds: explicit maps at their endpoints have multiple nearest affine maps. Below either threshold the distance is at most the inverse of the increasing small branch of Ψ . These conclusions are specific to the four-point test above. They are not statements about the identical threshold for every version of a three-quy homomorphism test.

A further consequence is the exact non-affine maximum

$$\mathcal{E}(f) \leq N^3 - 4N^2 + 10N - 6 \quad (N \geq 13),$$

with equality precisely for one-point modifications of affine maps (Corollary II.295.5). The cutoff thirteen is optimal: a parity-class example at $N = 12$ exceeds the displayed value. The last order requires a separate exact ternary quantization bound for thirteen character phases.

C. The exact interpolation list cost. Suppose an anchor a_i has a list of L_i base functions, while the fibre classes have degree less than t . A point is certified when it is sampled or its class contains t anchors. Source 39 uses

$$\sum_i L_i + E_t(L_1, \dots, L_r), \quad E_t(L_1, \dots, L_r) = \sum_{|I|=t} \prod_{i \in I} L_i \quad [269].$$

We remove the first term: the ordinary t -anchor interpolants also cover every sampled point. The answer is exactly E_t in the universal fixed-cell interface, with a separate complete formula when fewer than t lists are nonempty. In particular, for $r \geq 2$ and $L \geq 1$, the affine case improves $rL + \binom{r}{2}L^2$ to $\binom{r}{2}L^2$.

The exactness is not inferred from the upper count. We construct dense examples with one polynomial class per fibre that force all E_t candidates, even if the competing candidates have an arbitrary fixed degree in the base variable. For M candidates we give a product lower bound on the uncovered mass. This implies the asymptotically sharp lower bound $M \geq (1 - \delta - o(1))E_t$ when a fraction δ may be deleted, with integer rounding in the matching construction, and a cell-count/list-size tradeoff. The construction also shows exactly why a further compression must use information beyond this local interface.

D. Independent norm layers with no degree cap. Let q have characteristic $p > 2m$. Consider

$$C(x_1, \dots, x_s) = c + \sum_{j=1}^s v_j N_{\mathbb{F}_{q^{d_j}}/\mathbb{F}_q}(x_j), \quad v_j \in \mathbb{F}_q^r,$$

on the full direct product of its norm blocks. For $m = 2, 3, 4$, and for all $m \geq 2$ when $q \geq (2m - 1)^2$, the map avoids all nontrivial symmetrically coloured $2m$ -term progressions if and only if every d_j is odd and less than $2m$, and the coefficient vectors of each fixed degree are linearly independent (Theorem III.109.1). The exact maximum dimension is therefore rm^2 , without assuming any degree bound in advance.

The low-degree extremizing construction is due to Shi and Dong [33], and the repository already proves related degree-capped optimality statements [272]. The new part excludes every higher norm degree. Character sums give quantitative progression counts; exact positivity handles lengths four and six, and eleven explicit degree-nine irreducibility certificates handle the remaining eight-term fields. For general m , only degree $2m + 1$ over finitely many small prime fields remains outside the classification.

F.72.2 What is closed, and what remains open

Problem in the preceding work	Resolution and precise boundary
Unequal Fourier budgets near the common-budget branch	Exact formula, unique extremizer, fixed-certificate polytope, and Hessian rank. Farther phase regions remain open.
Range of the cubic affine-error profile	Entire positive range for exponent-three targets, with equality and stability. General odd-order targets remain a separate problem.
Universal ternary affine-recovery thresholds	Sharp strict thresholds $1/4$ and $4/9$ for the two specified domain classes. An exact envelope for every attainable defect is not asserted.

Problem in the preceding work	Resolution and precise boundary
Generic interpolation compression and lower bounds	Exact local cost, dense matching examples, and deletion/partition tradeoffs. Additional product-property information may still compress the lists.
Whether higher independent norm layers can improve the construction	Complete negative answer for lengths 4, 6, 8 and the stated large-field regime; an exact classification of all such layers.
Unrestricted scalar quintics for four-term symmetry	Still open. Interacting polynomial terms fall outside the independent norm-block theorem.

The distinction between a local theorem and a complete quantitative Szemerédi theorem matters. An improvement at one step must survive all subsequent losses, with compatible hypotheses at every scale, before it implies a new bound for the largest progression-free set $r_k(N)$. For comparison, Leng, Sah, and Sawhney prove a global bound of the form $r_k(N) \ll_k N \exp(-(\log \log N)^{c_k})$ for $k \geq 5$ and a positive c_k [2]. No result of that form is derived in this article. The contributions here are exact local theorems, equality classifications, and barriers that can be inserted into a further quantitative analysis.

F.72.3 Attribution, proof status, and organization

The main historical inputs are Gowers’ proof, the specified ProveIt reports, the supported Toeplitz moment criterion of Yang–Xie, elementary finite Fourier analysis, Lagrange interpolation, and the multiplicative Weil bound [1, 267, 152, 274, 275]. General stability of nearly extremal uniformity norms has a substantial prior literature; Eisner–Tao is one relevant reference [7]. We do not present near-extremizer stability itself as a new concept. Recent work of Mittal–Roy develops much broader low-soundness homomorphism tests using group-dependent query distributions [273]. Our thresholds concern the fixed uniform four-point graph-energy test, with explicit ternary equality cases.

The contribution statements are relative to a recorded repository and working-manuscript baseline. A targeted primary-source and repository audit was undertaken, including the more recent character/affine and bilinear-extraction manuscripts. It does not establish exhaustive historical priority for every corollary. The supplied source ledger records the inspected inputs without silently identifying a moving branch with an immutable snapshot.

Every asserted theorem has a proof below. The general results do not rest on random experiments. There is one explicit finite certificate component: the eleven irreducibility checks in the small-field eight-term norm argument. Their criterion, polynomials, and exact verifier are provided. Other exhaustive checks and numerical diagnostics serve to test calculations, endpoint cases, and normalizations. The manuscript received independent internal proof reviews during its preparation; it has not been externally refereed or formally verified in Lean.

Sections I.76–III.109 contain the four research developments. Section V.254 proposes further questions with explicit targets and obstructions. Section V.255 describes the verification package and a practical route to repository integration and formalization. The sections use local notation;

in particular the letter R in the Fourier section and r in the interpolation or norm sections refer to different parameters.

[write] **7 October 2026; sources 60–69.** *Identifications within this report* (checked at the write). [268], [269] and [270] are sources 37, 39 and 11 (Sections I.49, II.151, II.22); [272] names sources 10, 14, 22 and 38; [271], the “working manuscript” `gowers-character-affine-refinements.tex`, is source 59 as delivered (its SHA-256 in this source’s ledger equals that of source 59’s manuscript in arrival `9458a55a9`), and the ledger’s two further manuscripts are sources 49 and 57 as delivered, by the same check. The “merged report” it read is this article of the seventh write (blob `f4e5705`, sources 01–39); its README blob `cd43de1` is the seventh write’s README with one extra final line feed (Section F.4). *This source records no pin.* Its four suites were rerun on copies at the write.

F.73 Source 68. Introduction and contribution map

Source 68, delivered Section 1.

Added in the eleventh write, 7 October 2026. Delivered title block. “Sharp energy gaps and efficient restriction / Coset rigidity, puncture stability, and Fourier-filter bounds in Gowers’ Szemerédi framework / Research contribution prepared for Vladimir Reshetnikov / ProveIt / Combinatorics / Ramsey / Mathematical derivation and manuscript assisted by OpenAI ChatGPT / 7 October 2026”; *the PDF author field reads* “Research contribution prepared for Vladimir Reshetnikov with ChatGPT”. Archive `Gowers_Sharp_Energy_and_Restriction.zip` (arrival `0281fcb05`), manuscript `gowers_sharp_energy_restriction.tex` (3,422 lines, 44 pages), the standalone expansion of its `main.tex` and eleven section files by its `code/build_standalone.py` (shipped as `code/68-sharp-energy-build_standalone.py`; the expansion was reproduced byte for byte at the write), pin `0d93c1eac` (the arrival commit of sources 60–64; its `data/source_manifest.json`, shipped as `data/68-sharp-energy-source_manifest.json`, also records the archives of sources 45, 58 and 59 at `8ed496b14` and `f8bc5e2ec`). Printed below: the abstract, its keywords and Section 1; Sections 2–4 open the Part II chapter on the energy levels (Section II.296), Sections 5–7 open the Part II restriction chapter (Section II.310), Sections 8–9 are in Part V. Its macros `\sym` (the symmetric difference, not this report’s “SymMult”) and `\cost` (calligraphic $\mathcal{A}$, not this report’s $\mathcal{C}$) are set as this report’s `\syndiff`, the same definition, and as `\costA`; its ‘question’ environment (“Question”, on the shared counter) is set in this report’s ‘Question’ style; its figure is shipped as `figures/68-sharp-energy-overview.pdf`. The comments of its expansion inside sections are kept; those marking the joins between section files are dropped with the section breaks.

Abstract of source 68 (as delivered). We prove several local strengthenings of the additive-energy and restriction machinery surrounding Gowers’ proof of Szemerédi’s theorem. For a finite set A of size m in any abelian group, with energy defect $\varepsilon = 1 - E(A)/m^3$, we establish the sharp universal inequality $m r_A(d)(m - r_A(d)) \leq m^3 - E(A)$. An exact nonnegative remainder gives a mixed-set extension, a bounded weight version using minimum overlap, all one-translation equality cases, and a multilevel bound from the attained overlap spectrum. These yield a canonical subgroup and the sharp coset-distance envelope $|A \triangle C|/m \leq (1 - \sqrt{1 - 4\varepsilon})/2$ throughout $\varepsilon \leq 1/4$. The threshold $1/4$ is optimal even in finite cyclic groups; uniqueness of the closest coset has the distinct sharp threshold $2/9$. We also prove an exact nested-puncture repair profile with outer defect below $1/4$ and inner defect at most $1/4$, classify its equality cases, and obtain an oriented repair theorem from one nearly saturated translation.

For Gowers’ sixteen-variable restriction step, a degree-twelve rational filter with exact integer certificates improves the all-cyclic common exponent from the repository’s 54 to 53, with an explicit size hypothesis and arbitrary abelian targets. Finally, a circle rearrangement estimate and finite approximations to theta filters give bounds uniform in arity m and discrepancy order a . In particular the optimal principal-spectrum cost divided by m tends to $1 + 1/(2c)$ when $a = m^{c+o(1)}$. The manuscript supplies full proofs, sharp examples, reproducible certificates, a pinned source comparison, and a further research programme. Its claims concern the stated local

theorems and method cost; no improved global bound for Szemerédi’s theorem or literature-wide priority claim is made.

Keywords. Additive energy; stability; finite abelian groups; punctured cosets; Fourier restriction; Riesz–Sobolev inequality; Szemerédi’s theorem.

Gowers’ proof of Szemerédi’s theorem uses quantitative relations between additive configurations, approximate algebraic structure, and restriction of a map to a more additive domain [1]. Two local questions are especially amenable to sharp analysis. First, if a set has nearly the largest possible additive energy, how precisely is its structure forced, and what do almost sharp inequalities say about the missing points? Second, how efficiently can a random Fourier filter suppress unrespected additive relations while retaining respected ones? This article advances both questions and relates its conclusions to the existing ProveIt research sequence.

The first part is elementary finite additive combinatorics. Its basic identity keeps the remainder of a union–intersection transformation instead of discarding it. A single overlap gives a sharp quadratic gap; two overlap levels give a larger structural range. The resulting coset theorem has an optimal existence threshold and a different optimal uniqueness threshold. A second exact identity then reads the energy of the missing set inside the canonical coset, giving a nested repair theorem with a sharp nonlinear profile.

The second part keeps Gowers’ random-restriction architecture and the repository’s all-cyclic expectation theorem. We improve its explicit filter, and then study the extremal principal cost when both the number of variables and the order of a failed target relation vary. The finite coefficient improvement and the asymptotic theorem are separate results: their proofs use different filters.

F.73.1 Main statements in compact form

For a nonempty finite $A \subset G$, let $m = |A|$, $r_A(d) = |A \cap (A + d)|$, and $\varepsilon = 1 - E(A)/m^3$. Write $\tau(u) = (1 - \sqrt{1 - 4u})/2$. The first group of results is as follows.

1. **Universal gap and spectrum refinement.** For every translation d ,

$$\frac{r_A(d)}{m} \left(1 - \frac{r_A(d)}{m} \right) \leq \varepsilon.$$

More generally, if $g_1, \dots, g_\ell$ are the successive gaps between selected attained overlap values, together with 0 and m , then

$$E(A) \leq \frac{2m^3 + \sum_i g_i^3}{3}.$$

Theorems II.297.1 and II.297.4 prove these bounds. Theorem II.297.7 classifies all equality cases of the one-translation bound; Proposition II.297.8 supplies the minimum-overlap weighted extension.

2. **Coset existence and uniqueness.** If $\varepsilon \leq 1/4$, then $H = \{d : r_A(d) > m/2\}$ is a subgroup, and a most occupied H -coset C satisfies $|A \triangle C| \leq \tau(\varepsilon)m$. The threshold is optimal, including in cyclic groups. For $\varepsilon < 2/9$, C is the unique closest subgroup coset; the threshold $2/9$ is also optimal. See Theorem II.298.1, Corollary II.298.4, and Proposition II.298.6.

3. **Nested and oriented repair.** For $\varepsilon < 1/4$, use the canonical C and put $\delta = |A \triangle C|/m > 0$, $\eta = \varepsilon - \delta(1 - \delta)$. If $s = \eta/\delta^3 \leq 1/4$, then

$$\min_{D \subsetneq C} \frac{|A \triangle (C \setminus D)|}{m} \leq \delta\tau(s),$$

with D a nonempty subgroup coset. The analogous bound with η replaced by $\theta_d = \varepsilon - q(d)(1 - q(d))$ forces $d \in H$ and allows a minimizing hole whose subgroup excludes d . Theorems II.299.4–II.299.8 and Proposition II.299.11 include equality, endpoint, and uniqueness statements with their exact ranges.

For restriction, let B be a finite subset of a cyclic group, $b = |B|$, and $\phi : B \rightarrow F$ a map to an arbitrary abelian group. Let T_{16} count ordered tuples satisfying $x_1 + \cdots + x_8 = x_9 + \cdots + x_{16}$ and G_{16} count those also respected by ϕ . Repetitions are included. If $G_{16}(B, \phi) \geq \alpha b^{15}$, $0 < \alpha, \eta \leq 1$, and

$$b \geq \left(\frac{4}{\alpha\eta}\right)^{53},$$

Theorem II.312.2 gives one subset $B' \subset B$ with

$$T_{16}(B') \geq \left(\frac{\alpha\eta}{4}\right)^{53} b^{15}, \quad G_{16}(B', \phi) \geq (1 - \eta)T_{16}(B').$$

Theorem II.311.5 records and proves the predecessor expectation estimate used here. The change is a fully specified rational filter and the resulting stronger certified constants.

For the asymptotic question, let $\mathcal{I}_{m,a}$ be the infimum of $-\log P_m / \log(P_m/Q_{m,a})$ over finite even trigonometric polynomials $0 \leq K \leq 1$, $K(0) = 1$, with nonnegative Fourier coefficients c_n , where $P_m = \sum_n c_n^m$ and $Q_{m,a} = \sum_{a|n} c_n^m$. Theorem II.313.1 proves, for every $m, a \geq 2$,

$$\begin{aligned} \mathcal{I}_{m,a} &\geq \max \left\{ m - 1, m - 1 + \frac{m(\log m - 20)}{2 \log a} \right\}, \\ \mathcal{I}_{m,a} &\leq m - 1 + \frac{m [\log m + \log \log(e + \log m / \log a) + 4]}{2 \log a}. \end{aligned}$$

In particular, if $a = m^{c+o(1)}$ with fixed $c > 0$, then $\mathcal{I}_{m,a}/m \rightarrow 1 + 1/(2c)$. The additive term $m - 1$, negligible in the fixed-order leading asymptotic, becomes essential in this joint regime.

F.73.2 Exact comparison with the inspected repository

The repository changes rapidly. All current-tree comparisons in this article use the pinned commit

0d93c1eac2e20c313e7ed0cb7cb74ff705b00101.

The consolidated `article.tex` at the inspected state includes sources 01–39. Later placed material was checked separately; relevant full manuscripts for sources 45, 58, and 59 were retrieved from their historical incoming archives. Section V.257 records the paths and archive pins. Source numbers below are the repository’s identifiers, not publication numbers or proposed numbering for this contribution.




Figure F.73.1: Two proved profiles. Left: the canonical coset envelope, with punctured-coset equality parameters marked. The blue region has unique closest-coset guarantees; the remaining interval through $1/4$ retains existence. Right: the asymptotic principal cost when the discrepancy order is a power of the arity. This is a limiting curve, not a finite-arity numerical estimate.

Question	Inspected predecessor	Result here
One-translation energy gap	Source 45: $q(1 - q) \leq 2\varepsilon$; coefficient one asked	Sharp coefficient one, mixed and weighted forms, all set equality cases
Coset envelope	Source 14: $\varepsilon \leq 1/100$; source 45: $\varepsilon < 1/9$	Existence through $1/4$; uniqueness below $2/9$; both thresholds optimal in their stated senses
Nested-hole profile	Source 45: outer and inner defects below $1/9$	Outer defect below $1/4$, inner defect at most $1/4$; oriented version and full equality classification
All-cyclic, sixteen-variable restriction	Gowers: 2^{19} ; source 58: 54	Common exponent 53, explicit size condition, arbitrary abelian target
Principal filter cost	Source 58: fixed a , leading term $m \log m / (2 \log a)$ with $O_a(m \log \log m)$ error	Uniform two-parameter bounds; polynomial-order limit; fixed- a upper error $O_a(m(1 + \log \log \log m))$

The consolidated label `gsr:fz:q:II:cosetrange` asks for the coset range, exact envelope examples, and coordinate splitting. The present results answer its general-group range and canonical-equality parts; coordinate splitting and Cartesian structure are not resolved here. Source 45 separately asks the coefficient-one, largest-envelope-range, and nested-hole-range questions. Source 58 asks for a certified filter with $\mathcal{A}_{16,2} < 52$ and leaves the growing-order regime open. Theorems II.297.1, II.298.1, and II.313.1, and Proposition II.312.1, answer the specified coefficient-one, coset-range, cost-below-52, and leading growing-order questions. The nested result covers the full real domain of the inner quadratic inverse while keeping $\varepsilon < 1/4$. Proposition II.299.11 classifies the inner endpoint and shows why the outer endpoint fails.

F.73.3 Dependencies, interpretation, and limits

The energy part uses elementary finite counting, convolution, concavity, and subgroup arguments. The restriction part uses character orthogonality, Hölder’s inequality, and the already established decaying-error selection theorem from source 58, reproved here. The uniform cost lower bound additionally uses the classical Riesz–Sobolev rearrangement inequality on the circle [242]; the convolution concentration consequence needed here is derived explicitly. Poisson summation and finite Fejér approximation supply the upper bound.

These results refine specific parts of the Gowers framework. The near-maximal energy regime is stronger than the general large-energy hypothesis in the discussion preceding his Theorem 7.2. The exponent-53 theorem directly strengthens the quantitative conclusion of Lemma 9.3 under its displayed explicit size condition. Prime-field restriction results already in sources 04 and 17 have better retention powers in their own setting [277, 279]; the all-cyclic theorem is not a dominance claim over those results. The quantity $\mathcal{I}_{m,a}$ measures only principal-spectrum cost. It is not a lower bound on every possible subset-selection method, nor on every analysis of the same random subsets.

The report makes a targeted comparison with the inspected project and the cited primary analytic inputs. An exhaustive literature-priority determination is outside that comparison. In particular, standard union–intersection transformations, maximal-energy coset characterizations, character orthogonality, and random Fourier restriction are credited as established methods. All new mathematical assertions are the explicit statements proved below. No claim of a new global bound for progression-free sets or Ramsey numbers, a resolved longstanding major conjecture, or a new Lean kernel-checked theorem is made.

[write] **7 October 2026; sources 60–69.** *Identifications within this report.* [276] is this repository at 0d93c1eac, an ancestor of HEAD and the arrival commit of sources 60–64 (none of which it cites); the article it compared was the seventh write’s (blob f4e5705, sources 01–39), and its source manifest, shipped as `data/68-sharp-energy-source_manifest.json`, records the archives of sources 45, 58 and 59 by blob and SHA-256 at 8ed496b14 and f8bc5e2ec (verified at the write). Its sources 04, 14, 17, 45, 58 and 59 are this report’s; the label `gsr:fz:q:II:cosetrangle` it quotes is item II.31. *Not cited:* source 62, whose archive arrived in its pin commit and whose all-cyclic restriction has better powers for cyclic targets (Section II.310), and source 66, which arrived with it (Section II.296). Its four suites were rerun on copies at the write.

F.74 Source 69. The problem, the results, and the comparison baseline

Source 69, delivered Section 1.

Added in the eleventh write, 7 October 2026. Delivered title block. “Sharp All-Degree Boolean Phase Integration / Contraction radicals, extremal flags, and optimal finite-dimensional obstruction models / Research manuscript prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt project / 6 October 2026”; the PDF author field reads “Research manuscript prepared with ChatGPT for Vladimir Reshetnikov”. Archive `ramsey_boolean_phase_integration.zip` (arrival 0281fcb05), manuscript `article.tex` (1,636 lines, 25 pages), pin 0d93c1eac (its `SOURCES.md`, shipped as `69-obstruction-models-SOURCES.md`, records the blob of `Sections17_18.lean` there). Printed below: the abstract, its scope statement and Section 1; Sections 2–9 stand in the Part IV chapter on Boolean phase integration after source 65 (Section IV.51), Sections 10–14 and Appendix A are in Part V. Its macros `\F` ($\mathbb{F}_2$) and `\Torus` (the circle $\mathbb{R}/\mathbb{Z}$, not this report’s $\mathbb{T}$) are set as this report’s `\Ftwo`, the same definition, and as `\Tquot`; its ‘question’ environment (“Research question”, numbered 1–10 on its own counter) is set in this report’s ‘Research question’ style, numbered within its section, which keeps the indices.

Abstract of source 69 (as delivered). For a symmetric d -linear frequency tensor over $\mathbb{F}_2^n$, we determine exactly how much selected derivative Fourier energy is compatible with failure of nonclassical integration. The maximum is $1 - (d+1)/2^d$ for every $d \geq 3$ and every $n \geq 2$. This resolves the degree-seven question and the entire higher-degree gap left by the preceding project manuscript, whose universal answer was exact only through degree six. The new induction counts contractions that remain nonintegrable, rather than individual nonzero values of the repeated-variable defect. For $d \geq 4$ it also classifies all extremal tensor classes: they are indexed by flags consisting of a line inside a plane in the dual space, and are separated from all other nonintegrable classes by an explicit positive energy gap.

A second theorem gives the exact minimal quotient dimension of a tensor modulo integrable tensors. If its contraction-integrability radical has codimension ρ , this dimension is at most $\rho + \sum_{j=1}^{\min(\rho, d-2)} \binom{\rho}{j}$. An explicit construction attains this bound for every $d \geq 4$ and $\rho \geq 1$. The proof uses an integrable diagonal-lowering operator and a complete defect radical; it also gives a linear normal form for every prescribed radical. We give exact phase-removal and noisy-selector corollaries, algorithms for finite certificates, full proofs of the required background, and further research questions. Exact finite checks accompany the manuscript.

Scope and status. Unrefereed mathematical research draft. The claims of improvement are relative to the identified project baseline, not a certification of historical priority. No proof-assistant verification, unrestricted inverse theorem, or improved global bound for Szemerédi's theorem is claimed. Symmetry is a hypothesis, not a conclusion inferred from a general selector.

F.74.1 Selected derivative energy

Let $V = \mathbb{F}_2^n$. All averages in this manuscript are uniform normalized averages. For $f : V \rightarrow \mathbb{C}$ and $P : V \rightarrow \mathbb{R}/\mathbb{Z}$, set

$$\partial_h f(x) = f(x+h) \overline{f(x)}, \quad \Delta_h P(x) = P(x+h) - P(x), \quad e(t) = \exp(2\pi i t).$$

A function is *one-bounded* if $\|f\|_\infty \leq 1$; it need not have unit modulus and may vanish. Unless explicitly stated otherwise, all energy inequalities below concern one-bounded functions; exact algebraic identities may be applied to arbitrary complex functions. Write $\text{SM}_d(V)$ for the vector space of symmetric, genuinely d -linear maps $V^d \rightarrow \mathbb{F}_2$. This notation refers to multilinear maps, not to symmetric tensor coinvariants in characteristic two.

For $T \in \text{SM}_d(V)$ define

$$\mathcal{E}_T(f) = \mathbb{E}_{h_1, \dots, h_{d-1}} \left| \mathbb{E}_x \partial_{h_1} \cdots \partial_{h_{d-1}} f(x) (-1)^{T(h_1, \dots, h_{d-1}, x)} \right|^2. \quad (\text{F.74.1})$$

The last slot of T evaluates the selected Fourier frequency. In particular, $0 \leq \mathcal{E}_T(f) \leq 1$ for a one-bounded function. Let

$$\mathcal{M}(T) = \max_{\|f\|_\infty \leq 1} \mathcal{E}_T(f). \quad (\text{F.74.2})$$

The maximum exists because this is a continuous function on a finite product of compact complex discs.

A circle-valued function P has additive degree at most d when all its $(d+1)$ -fold additive differences vanish. Let $\mathcal{I}_d(V)$ be the space of *integrable* symmetric tensors: $T \in \mathcal{I}_d(V)$ when some such P has

$$\Delta_{h_1} \cdots \Delta_{h_d} P(x) = \frac{T(h_1, \dots, h_d)}{2} \pmod{1} \quad \text{for every } x, h_1, \dots, h_d. \quad (\text{F.74.3})$$

The notation $T/2$ embeds $\mathbb{F}_2$ as the two-torsion subgroup of $\mathbb{R}/\mathbb{Z}$. A primitive need not be an $\mathbb{F}_2$ -valued polynomial.

F.74.2 The exact universal threshold

Set

$$b_d = 1 - \frac{d+1}{2^d}. \quad (\text{F.74.4})$$

For independent $u, v \in V^*$ define the canonical obstruction

$$C_d^{u,v}(h_1, \dots, h_d) = \sum_{i=1}^d v(h_i) \prod_{j \neq i} u(h_j). \quad (\text{F.74.5})$$

Theorem F.74.1 (All-degree threshold and extremal classification). *For every $d \geq 3$ and $n \geq 2$,*

$$\max_{T \in \text{SM}_d(V) \setminus \mathcal{I}_d(V)} \mathcal{M}(T) = 1 - \frac{d+1}{2^d}. \quad (\text{F.74.6})$$

For $d \geq 4$, equality $\mathcal{M}(T) = b_d$ holds precisely for the tensors

$$T = C_d^{u,v} + S, \quad u, v \text{ independent}, \quad S \in \mathcal{I}_d(V). \quad (\text{F.74.7})$$

Every other nonintegrable symmetric tensor satisfies

$$\mathcal{M}(T) \leq 1 - \frac{3d}{2^{d+1}} = b_d - \frac{d-2}{2^{d+1}}. \quad (\text{F.74.8})$$

Consequently, energy strictly larger than b_d certifies integrability.

[write] **7 October 2026; sources 60–69.** *Second route* (Section IV.51): source 65's Theorem F.70.1 proves the same maximum, with the classification also at $d = 3$ and the radical-sensitive cap; the two inductions are the same in substance (integrable contractions form the radical $R(T)$, source 65's K_T), written independently.

The gap in (F.74.8) is a proved uniform separation, not a claim that the largest noncanonical energy always equals the displayed upper bound. The theorem classifies extremal *tensors modulo integrable tensors*; it does not classify all extremizing functions.

F.74.3 Optimal obstruction models

For $d \geq 3$ the repeated-variable defect is

$$\Phi_T(z_1, \dots, z_{d-3})(a, b) = T(a, a, b, z_1, \dots, z_{d-3}) + T(a, b, b, z_1, \dots, z_{d-3}). \quad (\text{F.74.9})$$

The frozen tuple is empty for $d = 3$. For $d \geq 4$, define

$$R(T) = \{z \in V : T_z \in \mathcal{I}_{d-1}(V)\}, \quad T_z(h_1, \dots, h_{d-1}) = T(z, h_1, \dots, h_{d-1}), \quad (\text{F.74.10})$$

$$\rho(T) = \text{codim}_V R(T). \quad (\text{F.74.11})$$

The subspace $R(T)$ is the radical of Φ_T in any one frozen slot. Thus ρ is a specific matrix-flattening rank. It is not the dimension of the span of the values of Φ_T , an alternating rank, or a partition rank.

Let $m_{\text{eff}}(T)$ be the smallest integer m for which

$$T = S + q^* T_0, \quad S \in \mathcal{I}_d(V), \quad q : V \twoheadrightarrow \mathbb{F}_2^m, \quad T_0 \in \text{SM}_d(\mathbb{F}_2^m). \quad (\text{F.74.12})$$

This measures the number of coordinates needed by the obstruction, after subtracting an integrable tensor.

Theorem F.74.2 (Sharp effective-dimension theorem). *For $d \geq 4$, put*

$$K(T) = \{a \in R(T) : \Phi_T(\mathbf{z})(a, b) = 0 \text{ for every } \mathbf{z} \in V^{d-3}, b \in V\}. \quad (\text{F.74.13})$$

Then

$$m_{\text{eff}}(T) = \text{codim}_V K(T) \leq B_d(\rho(T)), \quad B_d(r) = r + \sum_{j=1}^{\min(r, d-2)} \binom{r}{j}. \quad (\text{F.74.14})$$

For each $d \geq 4$ and $r \geq 1$, there is a symmetric tensor on $\mathbb{F}_2^{B_d(r)}$ with $\rho(T) = r$ and $K(T) = 0$. Thus the bound is sharp for every pair of parameters.

F.74.4 What is being improved

Gowers’s phase-removal step, especially Proposition 17.2 of [1], motivates the passage from structured derivative frequencies to polynomial phases. Tidor’s nonclassical integration theorem [150, Definition 3.1 and Proposition 3.5] supplies the correct small-characteristic notion of integrability. These are established inputs, not new claims of this manuscript.

The immediate comparison is the project manuscript *Exact Boolean Obstruction Energies* [284], source 47 of the inspected research collection. It already proves $\mathcal{M}(C_d^{u,v}) = b_d$ in every degree, the classification of pure defects, and the universal bounds

$$b_d \leq \max_{T \notin \mathcal{I}_d} \mathcal{M}(T) \leq 1 - \frac{\min(d+1, 7)}{2^d}. \quad (\text{F.74.15})$$

Its exact universal result therefore stops at degree six. Its first research question asks whether the degree-seven value is $15/16$, leaving the interval $120/128 \leq M \leq 121/128$. Theorem F.74.1 gives the exact value in that question and every higher degree at once.

Degree	Old universal upper bound	Exact value here	Improvement
4	11/16	11/16	already exact
5	13/16	13/16	already exact
6	57/64	57/64	already exact
7	121/128	$120/128 = 15/16$	1/128
8	249/256	247/256	2/256
9	505/512	502/512	3/512
$d \geq 7$	$1 - 7/2^d$	$1 - (d+1)/2^d$	$(d-6)/2^d$

The new ingredients are the contraction-integrability induction, its all-degree equality classification, and the sharp effective-dimension bound with extremizers. The canonical calculation, the integration criterion, and the cubic spectral argument are reproved and explicitly credited.

[write] **7 October 2026; sources 60–69.** *Identifications within this report.* [284] is source 47 (Section IV.42) and [283] source 43 (Section IV.25); [285] is this repository, read at `0d93c1eac`, where `Sections17_18.lean` has the blob `18495f5` recorded in its `SOURCES.md`, unchanged at HEAD `fb9602e55` (`proposition_17_2` at line 48). *Not cited, and not citing it:* source 65, which proves the main theorem independently on the same day; this source is printed after it as a credited second route (Section IV.51). The verifier was rerun on a copy at the write.

Part I

Density transfer and phase-flat partitions

Contents. The last step of a density increment: a correlation with a phase that is nearly constant on the cells of a partition into progressions is turned into a dense cell, and such partitions are built. Sections I.1–I.7: source 01, delivered Sections 2–8 (the finite weighted phase problem, the sharp density-tail theorem with occupied-mass exceptions, the optimal forced maximum, finite convex optimization, integer-exact simultaneous partitions, the Fourier increment, the scale barrier). Sections I.8–I.12: source 02, delivered Sections 6–9 and Appendix B (sharp real phase freezing, the sharp density–mass trade-off, the replacement for Lemma 5.15, the refinement theorem under the partition input $\text{PP}(K, R_0)$). Sections I.13–I.15: source 03, delivered Sections 2–4 (its conventions, the exact selection theorem, a density-sensitive Fourier increment). Sections I.16–I.20: source 05, delivered Sections 2–6 (the sharp constant $1/3$ in Lemma 2.3, Corollary 2.4 with 36π and $3\alpha/4$, the density-sensitive Fourier increment, simultaneous phases with the optimal exponent, spectral energy, short-cell discards).

Printed more than once, credited in dated notes. The sharp one-sided mass bound (sources 01, 02, 03, 05; note after Theorem I.4.1); the one-cell phase envelope (01, 02; note after Lemma I.1.3); the Lemma 5.15 replacement (02, 03); the simultaneous phase-flat partition (01, 05); the scale barrier (01 for one phase, 05 for m phases). The three Fourier increments (01, 03, 05) differ; source 01’s increment is at least $3/2$ times the others at the same length scale (note after Table 3).

Reading conventions of this Part (the full dictionary is Section F.3).

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
arc / phase width	01: an arc of r turns has half-angle $\theta = \pi r$ and $s = \sin \theta$; 02: chordal diameter 2ρ , so $\rho = s$; 03: chordal error $ e(\theta x) - z \leq \varepsilon$; 05: arc width ε turns, disk radius $\sin(\pi\varepsilon)$. <i>01’s r and 05’s ε are widths in turns, not angles.</i>
η, α, ρ	01: η = normalized correlation with $e(\xi n)$; 05: α = normalized Fourier coefficient; 03: ρ = normalized Fourier coefficient. For one coefficient 01’s η = 05’s α = 03’s ρ . 02: αN bounds a <i>sum over cells</i> of local correlations, and ρ is a phase half-diameter.
D, p, T, S, s, L	01: $D = \sum_j w_j q_j - \delta $ (twice the positive mass); 02: p positive mass, S sum of absolute cell sums, L total absolute mass; 03: $T = \mathbb{E}X_+$; 05: $S = \mathbb{E}g_+$, $s \leq S$, and $s(\varepsilon)$; 01: $s = \sin \theta$. <i>05’s s is a mass, 01’s s a sine.</i>
κ	01: $\kappa = \eta/(2\delta(1 - \delta))$; 05: a Gram-matrix bound (Corollary I.19.4).
Q, M, B, E	01: Q_i box counts, $M_\delta = 2\delta(1 - \delta)$, $B_{\delta,s}(a)$ largest correlation, E an exceptional set; 02: Q , $Q_{a,b}(p, t)$ good-mass certificates, M number of cells, B_0, B_ρ size constants; 05: $Q_i(L)$ box counts, M the modulus, $B = \mathbb{E} f $ and a width budget B , E captured Fourier energy, R an oscillation budget; 03: E_{osc} an oscillation error.

Added in the second write (6 October 2026): sources 06–10. Contents. Section I.21: source 06, delivered Section 5 (the exact trade-off between the mass of the cells above a threshold and the largest such cell for a fixed number M of cells, with exact extremizers; a strengthening of Lemma 5.15 with both branches of that optimum; the simultaneous improvement $\tau(\alpha)$; the indicator range; exceptional sets and signed drift). Sections I.22–I.24: source 10, delivered Sections 2–4 (its conventions; the optimal conversion of partition discrepancy into a density increment, its variance form and the mass above a threshold; correlation-to-density transfer, also with an exceptional set; an explicit phase partition with all rounding exposed and the one-parameter density–length frontier; the variance form of the U^2 inverse step).

Printed more than once (dated notes at both ends). The sharp one-sided mass bound is now proved in six sources: source 06’s Theorem I.21.1 (mass form, any total mass and signed mean) and source 10’s Proposition I.23.2 join sources 01, 02, 03 and 05 (note after Theorem I.4.1); the forced maximum is also source 10’s (I.23.2). The Lemma 5.15 replacement of sources 02 and 03 is sharpened by source 06’s fixed- M optimum (notes after Corollary I.9.3 and Corollary I.21.2). Exceptional mass appears in three different

models: source 01's Theorem I.2.3, source 06's Proposition I.21.4, source 10's Theorem I.23.4. Source 10's Corollary I.24.3 has source 01's length formula and source 05's increment; its frontier, Theorem I.24.2, is compared with source 01's increment in the note after Table 3.

Reading conventions for the sections of sources 06–10 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
$\eta, \tau, \gamma, \lambda$	06: η a density <i>threshold</i> in Theorem I.21.1, $\tau(\alpha)$ the simultaneous threshold of Corollary I.21.3. 10: η a normalized correlation (= 01's η , 05's α), γ the retained fraction of the correlation, τ a density threshold (= 06's η), λ a phase angle. <i>06's η is not a correlation; 10's λ is not 03's loss parameter.</i>
v, S, V, M, W	06: $v = \sum_j w_j x_j $ the cell variation (for $f = 1_A - \delta$ this is 01's D), m the signed mean, $P_{\pm} = (v \pm m)/2$, S the high-mass bound, W the total mass, M the number of cells, $[-a, b]$ the range (as 02, 03). 10: $S_{\mathcal{P}}$ the same first moment as 06's v , $V_{\mathcal{P}}$ the second moment, $v = \delta(1 - \delta)$, M the <i>largest cell density</i> , D_{γ} an increment. <i>10's v is not 06's v, and 10's M is not a number of cells.</i>
ρ, s, ε	06: ρ the discarded (exceptional) mass, ε a bound on the signed mean. 10: ρ the disc-approximation error of the phase, $s = \eta - 2v\rho$ the residual correlation, ε the exceptional mass. <i>10's s is not 01's $\sin \theta$.</i>

Added in the third write (6 October 2026): sources 11–14. Contents. Sections I.25–I.29: source 13, delivered Sections 2–6 (its conventions; simultaneous linear phases: complete partitions with flexible lengths, the optimal exponent frontier with a random-slope obstruction and a finite certificate, complete partitions with prescribed lengths $v - 1$ or v and exponent $1/(3q + 2)$, bounded-height presentations; the sharp high-average mass bound with discarded mass and drift; the real phase envelope and a sharp joint phase–density–exception frontier; a density-sensitive Fourier increment).

Printed more than once (dated notes at both ends). The sharp one-sided mass bound is now proved in seven sources (source 13's Theorem I.27.1); the one-cell envelope and the linear law (01, 02, 13); the joint frontier (source 01's Theorems I.2.1–I.3.1, source 13's Theorem I.28.3); the simultaneous partition (01, 05, 13); the obstruction (05's minimax theorem, 01's barrier, 13's random slopes with independent tolerances). Source 13's Fourier increment is dominated by source 01's at the same length (proved at the write, note after Theorem I.29.1). New in source 13: prescribed lengths with exponent $1/(3q + 2)$ and no size threshold, and independent tolerances in the obstruction.

Reading conventions for the sections of sources 11–14 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
q, Q_i, D, d	13: q the number of linear phases (05's m), Q_i approximation denominators, $D = \prod Q_i$, d a common difference, $a_q = 1/(3q + 2)$, $E_0 = \prod \varepsilon_i$, L_* (= 01's L), $B_v = (v - 1)(v - 2)$. <i>13's q is not a common difference; its B_v is not a translation budget.</i>
$D, w_j, x_j, P_{\pm}, R_p(t)$	13, Section I.27: $D = \sum w_j x_j $ (= 01's D , twice the positive mass), $R_p(t)$ (= 02's $Q_{a,b}(p, t)$), M the number of cells (= 02's and 06's M).
c, d, v, s, C, V	13, Section I.28: c, d, v the unnormalized correlation, signed and absolute integrals of one cell (<i>this d is not a difference</i>), $2s$ a chord diameter ($s = 02$'s $\rho = 01$'s s), C the aggregate correlation, $V = \mathbb{E} f $, $v(x)$, h_s , H_s , K_s the frontier functions.
α, ρ_0	13, Section I.29: α a normalized Fourier coefficient (= 01's $\eta = 05$'s α), $\rho_0 = \alpha/(4\delta(1 - \delta))$.

Added in the fourth write (6 October 2026): sources 15–20. Contents. Sections I.30–I.31: source 16, delivered Sections 6–7 (exact partitions whose own steps control affine frequencies: a certificate theorem,

the cited quadratic recurrence theorem, the power-law region with every balanced $\gamma < 1/((q+1)(42q+1))$, composite moduli, a weighted cell for the Lemma 13.5 interface; an endpoint-cell obstruction $\gamma \leq 1/(3q+1)$. Sections I.32–I.33: source 18, delivered Sections 4–5 (simultaneous polynomial partitions into lengths $v-1$ or v by descent in degree; the quadratic and shared-leading-coefficient exponents; a volume obstruction for every degree). Section I.34: source 19, delivered Section 6 (degree-profile partitions from Maynard’s theorem, target lengths, a necessary frontier for arbitrary degrees, finite diagnostics and five questions).

Overlaps (dated notes). Sources 18 and 19 prove the same recurrence-to-partition transfer independently; source 18’s prescribed-length exponent is the better one, source 19’s degree profiles the more general; their obstructions agree at uniform degree, and in degree one they are source 13’s results (Corollary I.26.9, Theorem I.26.5) and source 05’s. Source 16’s own-step problem is different (affine frequencies times the actual step, Lemma 13.5), with the same recurrence input. **Cross-source answers.** Source 13’s Question 3 is answered by sources 18 and 19, conditionally on Maynard’s published theorem (item I.1); source 13’s Question 1 is not answered by them (the placement of source 23 records an answer, for a later write).

Reading conventions for the sections of sources 15–20 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
own steps	16: q the number of affine frequencies $A_i n + B_i \in \mathbb{T}$ (13’s q ; <i>not a common difference</i>), δ_i tolerances on $(A_i n + B_i)d_Q$ with d_Q the cell’s actual step, $J_i = \lceil 2/\delta_i \rceil$, $T = \prod_i J_i$ (<i>not a group</i>), L a minimum length, s a certificate, $C > 21$ the recurrence constant, $X_0(C, q)$ its threshold, γ a balanced exponent, R an interval length.
partitions	18: k the degree, q the number of phases, v a target length, $g(v) = \max\{1, (v-1)(v-2)\}$, $\mathcal{P}_{k,q}(v, \varepsilon)$ a threshold, A_k, B_k, D_k exponent constants (<i>not 01’s or 04’s A_k, B_k</i>), osc the length of a containing arc, $\kappa > 10.5$ (Lau), s a scalar recurrence exponent, $H = k(k+1)/2$, $\mathbb{R}/\mathbb{Z}$ the circle. 19: q_j the number of phases of degree at least j , $\kappa_j, F_j, G_j, S = \sum_i d_i(d_i+1)/2, D = \sum_i d_i$, ad an arc diameter, θ a target-length exponent.

Added in the fifth write (6 October 2026): sources 21–27. Contents. Sections I.35–I.39: source 23, delivered Sections 2–6 (the partition problem and the conductor of the semigroup of allowed lengths; matching upper and lower thresholds; sharp asymptotics for lengths $L, L+1$ and for intervals of lengths; consequences for Lemma 5.9 and finite witnesses; exact packing and almost-partitions). Sections I.40–I.44: the chapter on Corollary 5.8 without a scale threshold, a guide and source 26’s delivered Sections 2–4 and Appendix A (threshold-free bounds; a moment increment over every cyclic modulus with removal of nonproper steps; phase-flat refinement; integer-box almost-covers).

Overlaps (dated notes). Source 23’s upper threshold is source 13’s Theorem I.26.8 in a general form (source 23 says so); its lower bound is new. Source 26’s Dirichlet step and phase-flat cells are of the kind of sources 01, 05 and 13, its positive-cell extraction that of sources 02, 03 and 06, at a different interface. **Cross-source answers.** Source 13’s Question 1 is answered by source 23 (item I.12).

Reading conventions for the sections of sources 21–27 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
thresholds	23: $I_R = \{0, \dots, R-1\}$, adiam an arc diameter (Gowers's convention, normalized; <i>not the largest pairwise distance</i>), q the number of affine phases (13's q), $\mathcal{S}$ the allowed lengths, $L = \min \mathcal{S}$, $H = \max \mathcal{S}$ (<i>not a group or a statistic</i>), $C(\mathcal{S})$ the conductor, $T_{\mathcal{S}}(\varepsilon)$ the eventual universal threshold, M_i radices, D a mixed radix, B the Dirichlet bound, $g_{\mathcal{S}}$, $A_{\mathcal{S}}$ gap functions.
Corollary 5.8	26: $G = \mathbb{Z}/N\mathbb{Z}$, proper progressions are injectively parameterized and may wrap (<i>not integer progressions</i>), U the union of the parents, $\tau = U /N$, r the shortest parent, μ the normalized balanced mass, α the summed correlation, L a progression length, Q the output progression, D_{bad} nonproper steps, $\text{spf}(N)$ the least prime factor, D a box dimension in Appendix A.

Added in the sixth write (6 October 2026): sources 28–33. Contents. Sections I.45–I.48: source 32's delivered Sections 2–5, which continue the chapter on Corollary 5.8 (an explicit quadratic recurrence; lossless phase partitions; a second moment over every cyclic modulus; assembly at every scale).

Overlaps (dated notes). Source 32's variance section is source 26's (the sibling reports share it). **Cross-source answers.** The degree-two part of source 26's Question 4 is answered by source 32 (item I.19).

Reading conventions for the sections of sources 28–33 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Corollary 5.8	32: $e(t) = \exp(2\pi it)$, $\ t\ $ the circle distance, proper progressions injectively parameterized (may wrap), phases $a_i x^2 + b_i x + c_i$ with coefficients in G (<i>not phases with vanishing third differences</i>), T a Weyl length (<i>not an energy</i>), R a recurrence radius, p the recurrence integer (<i>not a prime</i>), $H = 4R^2$ or $64L^3$ (<i>not a subgroup</i>), q a Dirichlet denominator, L the cell length, $S = N/M$, τ , r , μ , $b = 1 - \delta$ as in source 26.

Added in the seventh write (6 October 2026): sources 34–39. Contents. Sections I.49–I.59: a guide and source 37's delivered Sections 2–9 and Appendices A–B (the exact gap–moment theorem; a positive-coefficient polynomial certificate; a supported trigonometric moment criterion; attainment, support and equality; endpoints and the first spectral transition; a linear-frequency replacement for Lemma 5.2; stability; a stronger monomial-recurrence deduction; the real interval factorization; low-dimensional examples). Sections I.60–I.61: a guide and source 35's Section 7 (polynomial divisibility and prescribed progression lengths).

Overlaps (dated notes). Source 35's frontier is source 23's at degree one and adds the conductor to source 19's flexible-length obstruction for prescribed lengths. **Cross-source answers.** Source 23's Question 6 is answered in its construction half by source 35 (item I.15); no earlier source treats Lemma 5.2.

Reading conventions for the sections of sources 34–39 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Lemma 5.2	37: $\mathbb{T}_{2\pi} = \mathbb{R}/2\pi\mathbb{Z}$, I_a the open arc of radius a , K_a its complement, $\mathcal{P}(K_a)$ probability measures on K_a , $E_R(\mu) = \max_{r \leq R} \hat{\mu}(r) $, $m = R + 1$ (<i>not a modulus</i>), S_R , d_R , $\rho_R(a)$, $\mu_{R,a}$ the extremizer, w the certificate vector, M the gap half-width and N the modulus in Section I.55, D the Weyl exponent, $B = CT^\eta$ its constant, A_k , a_k , b_k recurrence constants and exponents, $\ \cdot\ _{\mathbb{R}/\mathbb{Z}}$ the circle distance.
Lemma 5.9	35: $d_1, \dots, d_\nu$ degrees, $W = \sum_i d_i(d_i + 1)/2$ ($= 19$'s S ; <i>not 35's $W = \sigma^\perp$ of Part II</i>), D the forced divisor (<i>not 19's $D = \sum_i d_i$</i>), $C(\mathcal{S})$ the conductor of the length set $\mathcal{S}$, $\mathcal{T}$ the threshold, L the minimum cell length, ε_i tolerances, a , b_i exponents.

Added in the eighth write (7 October 2026): sources 40–46. Contents. Sections I.62–I.74: a guide, source 41’s Sections 2–7 (an explicit quadratic recurrence estimate; explicit recurrence in every degree; lossless phase partitions; a second moment over every cyclic modulus; assembly at every scale; exact budgets and sharper alternatives) and source 42’s Sections 2–7 (the covariance square and the range inequality; exact support optimization and integer scales; positive definite weights and the classical packing bound; exact extremizers and the optimal uniform constant; a density-sensitive gain at the quadratic endpoint; a sparse affine cover consequence).

Overlaps (dated notes). Source 41’s Sections 2 and 5 are source 32’s Sections 2 and 4 verbatim, and its retained affine and quadratic bounds are sources 26’s and 32’s theorems; source 42’s variance step replaces the one of sources 26, 32 and 41. **Cross-source answers.** Source 32’s degree-three question and source 26’s higher-degree question are answered by source 41 (item I.19); source 26’s question on the $(N/8)^{1/3}$ ceiling is answered by source 42.

Reading conventions for the sections of sources 40–46 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Corollary 5.8	41: k the degree, $s = k - 1$, $w = 2^s$, K_k , D_k , A_k , B_k , m_k , c_k , e_k , $F_k(L) = 2^{c_k} L^{e_k}$, $\mathcal{P}_k$, R a recurrence accuracy, L the target length, H , T integer scales, $S = N/M$, $b = 1 - \delta$, $\tau = U /N$, $r = \min_i P_i $, μ the balanced mass, E a phase error (<i>not an energy</i>), λ , γ_* limiting constants. 42: m , q (<i>not a field size</i>), $M_L(A)$, $H_L(A)$, $S = B $ a support size (<i>not 41’s $S = N/M$</i>), β the proportion of short-order directions (<i>not a density</i>), η a collision probability, C a scale constant.

Added in the eleventh write (7 October 2026): sources 60–69. Contents. Sections I.75–I.76: a guide and source 67’s Section 2 (exact Fourier-gap profiles with unequal budgets), which continues source 37’s chapter on Lemma 5.2 (Section I.49).

Reading conventions for the sections of sources 60–69 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Lemma 5.2	67: I_a the open arc of radius a , $V_{R,a}(\varepsilon)$ the least mass of I_a under budgets $ \widehat{\mu}(j) \leq \varepsilon_j$, ρ source 37’s critical value ($= 37$ ’s $\rho_R(a)$), $L(\varepsilon)$ an $R \times R$ localizing matrix (<i>not a length</i>), $\mathbf{e}$ the all-ones vector, $c = \cos a$, Z the extremal support, f_j the certificate coefficients, $\rho_{R,N}$ the grid value.

I.1 The finite weighted phase problem

Source 01, delivered Section 2.

I.1.1 Conventions

Let (Ω, μ) be a finite probability space with positive atom weights. Let $h : \Omega \rightarrow [0, 1]$, and put

$$\delta = \mathbb{E}h \in (0, 1), \quad b = 1 - \delta, \quad f = h - \delta, \quad M_\delta = 2\delta b. \quad (\text{I.1.1})$$

A partition $\mathcal{P} = (P_j)$ consists of nonempty cells. Set

$$w_j = \mu(P_j), \quad q_j = \mathbb{E}(h \mid P_j).$$

Then $\sum_j w_j = 1$ and $\sum_j w_j q_j = \delta$. Every assertion for this model includes ordinary subsets of a finite set, by taking uniform weights and $h = \mathbf{1}_A$.

A unit-modulus function ψ is *short-arc on a cell* if there is a real γ_j such that

$$\psi(x) \in \{\exp(i(\gamma_j + t)) : -\theta \leq t \leq \theta\} \quad (x \in P_j), \quad 0 \leq \theta < \pi/2. \quad (\text{I.1.2})$$

Write $s = \sin \theta \in [0, 1)$. Arc centers may differ between cells. We assume

$$|\mathbb{E} f \psi| \geq \eta \geq 0. \quad (\text{I.1.3})$$

The elementary chord inequality

$$|u - \delta| \leq \delta + (1 - 2\delta)u \quad (0 \leq u \leq 1) \quad (\text{I.1.4})$$

gives $\mathbb{E}|f| \leq M_\delta$, and hence $\eta \leq M_\delta$. Equality holds in the first inequality for indicator functions.

I.1.2 A cancellation-sensitive preliminary bound

Lemma I.1.1 (Cancellation and phase diameter). *Let f be any real function on a finite probability space, and suppose $|\psi| \leq 1$ and the diameter of $\psi(P_j)$ is at most $2s$, where $0 \leq s \leq 1$. If*

$$M = \mathbb{E}|f|, \quad D = \sum_j \left| \mathbb{E}(f \mathbf{1}_{P_j}) \right|,$$

then

$$|\mathbb{E} f \psi| \leq sM + (1 - s)D. \quad (\text{I.1.5})$$

Proof. In one cell let A and B be the positive and negative masses of f , measured with the original probability weights. Choose the corresponding phase averages u, v (arbitrarily within the phase hull if a mass is zero). They belong to the convex hull of the cell's phase values, so $|u|, |v| \leq 1$ and $|u - v| \leq 2s$. If $A \geq B$, then

$$|Au - Bv| = |(A - B)u + B(u - v)| \leq A - B + 2sB.$$

This equals $s(A + B) + (1 - s)|A - B|$. Interchanging the two signs proves the same estimate when $B > A$. Sum over cells and apply the triangle inequality. $\square$

For indicators, the coefficient M in (I.1.5) is $2\delta(1 - \delta)$, not the cruder bound 1. Consequently,

$$D \geq \frac{\eta - 2\delta(1 - \delta)s}{1 - s} \quad (\text{I.1.6})$$

whenever the right-hand side is positive. The next calculation retains more information than D alone.

I.1.3 The exact one-cell envelope

Define, for $0 \leq q \leq 1$,

$$\Phi_{\delta,s}(q) = \sqrt{(q - \delta)^2 + 4\delta b q(1 - q)s^2}. \quad (\text{I.1.7})$$

Lemma I.1.2 (Sharp cell bound). *Under (I.1.2),*

$$|\mathbb{E}((h - \delta)\psi \mid P_j)| \leq \Phi_{\delta,s}(q_j). \quad (\text{I.1.8})$$

For each $q \in [0, 1]$, equality can be realized on a finite weighted cell with mean q and an indicator h .

Proof. First suppose h is an indicator. Relative to the cell probability measure, the positive mass is $A = bq$ and the negative mass is $B = \delta(1 - q)$. Let U, V be the phase averages on the occupied and unoccupied portions. Taking independent samples of their phase values,

$$|AU - BV| \leq \max_{|u|, |v| \leq \theta} |Ae^{iu} - Be^{iv}|.$$

The omitted common rotation has modulus one. Since $|u - v| \leq 2\theta < \pi$,

$$|Ae^{iu} - Be^{iv}|^2 = (A - B)^2 + 4AB \sin^2 \frac{u - v}{2} \leq (A - B)^2 + 4ABs^2.$$

Here $A - B = q - \delta$, giving (I.1.8). Place all occupied mass at one endpoint of the arc and all unoccupied mass at the other to attain it. Cases $q = 0, 1$ use only one phase value.

For general h , replace an atom of weight w and value $h(x)$ by an occupied atom of weight $wh(x)$ and an unoccupied atom of weight $w(1 - h(x))$, with the same phase. Delete zero-weight atoms. This preserves both the cell mean and the correlation, so the indicator argument applies. $\square$

Lemma I.1.3 (Convexity and the endpoint chord). *The function $\Phi_{\delta,s}$ is convex on $[0, 1]$, with $\Phi_{\delta,s}(0) = \delta$ and $\Phi_{\delta,s}(1) = b$. Moreover,*

$$\Phi_{\delta,s}(q) < \delta + (1 - 2\delta)q \quad (0 < q < 1, \ 0 \leq s < 1). \quad (\text{I.1.9})$$

Proof. For $s = 0$, this is the convex function $|q - \delta|$. For $0 < s < 1$, differentiation gives

$$\Phi''_{\delta,s}(q) = \frac{4\delta^2 b^2 s^2 (1 - s^2)}{\Phi_{\delta,s}(q)^3} > 0. \quad (\text{I.1.10})$$

For completeness, if $G = \Phi_{\delta,s}^2$ then $(\sqrt{G})'' = (2GG'' - (G')^2)/(4G^{3/2})$, and the numerator divided by four is $4\delta^2 b^2 s^2 (1 - s^2)$. The endpoint values are immediate. Finally, the square of the right-hand side of (I.1.9), minus $\Phi_{\delta,s}(q)^2$, is $4\delta b q(1 - q)(1 - s^2) > 0$. The chord is positive, so taking square roots proves the strict inequality. $\square$

[write] 6 October 2026. *One bound, two parametrizations.* Lemma I.1.2 and source 02's Theorem I.8.1 are the same bound. In a cell with positive mass $A = bq$ and negative mass $B = \delta(1 - q)$ (cell-normalized), source 02's signed and absolute masses are $F = A - B = q - \delta$ and $L = A + B$, and its bound $C^2 \leq (1 - \rho^2)F^2 + \rho^2 L^2 = F^2 + 4\rho^2 AB$ is $\Phi_{\delta,s}(q)^2$ with $\rho = s$. The arc conventions agree: an arc of half-angle θ has chordal diameter $2 \sin \theta$, so source 02's ρ (half the chordal diameter) is this $s = \sin \theta$. Both proofs sample the phases of the positive and the negative mass independently; source 02's statement is for arbitrary real f , and this one reduces bounded weights to indicators by splitting atoms. The linear consequence, Lemma I.1.1, is source 02's (I.8.2) and Corollary I.8.2, proved here directly and there from the quadratic bound. Credited to both sources. Source 05's Proposition I.17.2 (a disk centre inside the arc) and source 03's Lemma I.14.3 (arbitrary unit constants) are different, complex-valued transfer estimates.

[write] 6 October 2026; sources 11–14. Source 13 proves the envelope and the linear law a third time (its Theorem I.28.1 and Corollary I.28.2), and the joint frontier of Theorems I.2.1–I.3.1 as its Theorem I.28.3 (notes there).

I.2 A sharp density-tail theorem

Source 01, delivered Section 3.

Fix a target density $a \in [\delta, 1)$. Define

$$u_a = \frac{\Phi_{\delta,s}(a) - \delta}{a}, \quad (\text{I.2.1})$$

$$B_{\delta,s}(a) = \delta + \delta u_a, \quad (\text{I.2.2})$$

$$\Gamma_{\delta,s}(a) = 1 - 2\delta - u_a = \frac{M_\delta - B_{\delta,s}(a)}{\delta} > 0. \quad (\text{I.2.3})$$

Positivity in (I.2.3) follows from (I.1.9). The quantity $B_{\delta,s}(a)$ is the largest correlation permitted by these constraints if every cell density is at most a .

Theorem I.2.1 (Sharp excess and density-tail transfer). *Under the hypotheses of Section I.1, put*

$$\lambda_*(\delta, s, a, \eta) = \frac{(\eta - B_{\delta,s}(a))_+}{\Gamma_{\delta,s}(a)} = \frac{\delta (\eta - B_{\delta,s}(a))_+}{M_\delta - B_{\delta,s}(a)}. \quad (\text{I.2.4})$$

Then

$$\sum_{q_j > a} w_j \geq \sum_j w_j \frac{(q_j - a)_+}{1 - a} \geq \lambda_*(\delta, s, a, \eta). \quad (\text{I.2.5})$$

For every $\eta \in [0, M_\delta]$, the last bound and the corresponding lower bound on the mass of cells with $q_j > a$ are optimal among all finite weighted spaces and phases satisfying the stated constraints.

Proof. Convexity bounds $\Phi_{\delta,s}$ by the chord joining 0 to a on the first interval, and by the chord joining a to 1 on the second. Written as one formula, this is

$$\Phi_{\delta,s}(q) \leq \delta + u_a q + \Gamma_{\delta,s}(a) \frac{(q - a)_+}{1 - a}. \quad (\text{I.2.6})$$

To check the second chord, its slope minus u_a is

$$\frac{b - \Phi_{\delta,s}(a)}{1 - a} - u_a = \frac{b - \delta - u_a}{1 - a} = \frac{\Gamma_{\delta,s}(a)}{1 - a}.$$

Lemmas I.1.2 and I.1.3, the triangle inequality, and $\sum_j w_j q_j = \delta$ now give

$$\eta \leq \sum_j w_j \Phi_{\delta,s}(q_j) \leq B_{\delta,s}(a) + \Gamma_{\delta,s}(a) \sum_j w_j \frac{(q_j - a)_+}{1 - a}. \quad (\text{I.2.7})$$

Rearrange and use $0 \leq (q - a)_+/(1 - a) \leq \mathbf{1}_{q > a}$.

For sharpness, first take $\eta \geq B_{\delta,s}(a)$ and write $\lambda = (\eta - B_{\delta,s}(a))/\Gamma_{\delta,s}(a)$. Since $\eta \leq M_\delta$, we have $0 \leq \lambda \leq \delta$. Give cells of densities 1, a , 0 respective weights

$$w_1 = \lambda, \quad w_a = \frac{\delta - \lambda}{a}, \quad w_0 = 1 - \lambda - \frac{\delta - \lambda}{a}. \quad (\text{I.2.8})$$

They are nonnegative, sum to one, and have weighted mean δ . Use the equality construction of Lemma I.1.2 in each cell, and rotate the arcs independently so that all cell correlations point along the positive real axis. The resulting correlation is

$$w_0 \delta + w_a \Phi_{\delta,s}(a) + w_1 b = B_{\delta,s}(a) + \Gamma_{\delta,s}(a) \lambda = \eta.$$

Only the density-1 cells exceed a , and their normalized excess is also λ . Thus both bounds in (I.2.5) are equalities.

For $0 \leq \eta \leq B_{\delta,s}(a)$, use the $\lambda = 0$ configuration. If $B_{\delta,s}(a) > 0$, take two disjoint copies with opposite global phase rotations and relative weights $(1 + \eta/B_{\delta,s}(a))/2$ and $(1 - \eta/B_{\delta,s}(a))/2$. They still have mean δ , no cell exceeds a , and their total correlation is η . If $B_{\delta,s}(a) = 0$, only $\eta = 0$ arises in this case. Zero-weight cells are omitted throughout. $\square$

Remark I.2.2 (The scope of sharpness). The extremizers allow arbitrary real atom weights and independent choices of arc centers. With rational weights they can be converted to ordinary finite sets by replacing an atom with equally weighted copies. General weights can be approximated by rational ones. This is not a proof that the same extremizers occur for a fixed character on $\mathbb{Z}/N\mathbb{Z}$, a fixed progression partition, or a prescribed list of cell sizes. Such arithmetic constraints may permit further improvements.

I.2.1 Exceptional cells cost occupied mass

A usual exceptional-set estimate pays for the whole exceptional set. Here its occupied mass is the relevant quantity.

Theorem I.2.3 (Occupied-mass robust transfer). *Let $E \subseteq \Omega$, and partition $\Omega \setminus E$ into nonempty regular cells P_j . Impose the short-arc condition only on regular cells, and set*

$$\beta_E = \mathbb{E}(h\mathbf{1}_E).$$

Retain the global mean δ and correlation lower bound η . For $\delta \leq a < 1$,

$$\sum_{\substack{j \text{ regular} \\ q_j > a}} w_j \geq \sum_{j \text{ regular}} w_j \frac{(q_j - a)_+}{1 - a} \geq \max\{0, \lambda_*(\delta, s, a, \eta) - \beta_E\}. \quad (\text{I.2.9})$$

Here $w_j = \mu(P_j)$ uses the original measure, not the measure conditioned on the regular set.

Proof. Write $\varepsilon = \mu(E)$. On regular cells use (I.2.6); their total weight and occupied mass are $1 - \varepsilon$ and $\delta - \beta_E$. On E , use (I.1.4). Thus

$$\begin{aligned} \eta &\leq \delta(1 - \varepsilon) + u_a(\delta - \beta_E) + \Gamma_{\delta,s}(a) \sum_{j \text{ regular}} w_j \frac{(q_j - a)_+}{1 - a} + \delta\varepsilon + (1 - 2\delta)\beta_E \\ &= B_{\delta,s}(a) + \Gamma_{\delta,s}(a) \left(\beta_E + \sum_{j \text{ regular}} w_j \frac{(q_j - a)_+}{1 - a} \right). \end{aligned}$$

Rearranging proves the result, including the case $\eta \leq B_{\delta,s}(a)$. $\square$

The estimate is especially useful when the excluded cells have little of A , even if their total length is not particularly small. An exceptional set consisting of unoccupied points has $\beta_E = 0$: it incurs no loss in this formula. There is no contradiction with global cancellation, since the density and constant term were never renormalized after deletion.

[write] **6 October 2026; sources 06–10.** *Exceptional mass in three models.* Source 06's Proposition I.21.4 (extraction from a cell variation after a set of mass ρ is discarded, or with a signed mean of size at most ε) and

source 10's Theorem I.23.4 (correlation transfer with an exceptional set of mass $\varepsilon < \delta$, under $s > 2(1 - \delta)\varepsilon$) also allow an exceptional set. The theorem above charges only the occupied mass $\beta_E \leq \mu(E)$; the other two charge the total exceptional mass. The three statements have different hypotheses and conclusions, and none is claimed to imply another.

Corollary I.2.4 (A large good cell). *For a uniform space of N points with m regular cells, let $\rho = \max\{0, \lambda_* - \beta_E\}$. If $\rho > 0$, some regular cell has*

$$q_j > a, \quad |P_j| \geq \frac{\rho N}{m}.$$

If all regular cells have size at least L , that cell also has size at least L . If instead all have size at most U , there are at least $\rho N/U$ good cells.

Proof. Their union has at least ρN points by Theorem I.2.3. Pigeonhole over at most m cells for the first assertion; divide their total size by the upper size bound for the last. $\square$

I.3 The optimal forced maximum and its tradeoff

Source 01, delivered Section 4.

The density-tail theorem contains a closed-form answer to a simpler question: how large must at least one cell density be?

Theorem I.3.1 (Optimal maximum density from a short-arc correlation). *Assume there are no exceptional points, and put $\kappa = \eta/M_\delta$. If $0 \leq s < \kappa \leq 1$, then*

$$\max_j q_j \geq a_* = \frac{\delta(\kappa - s^2)}{\kappa(1 - b\kappa) - \delta s^2}. \quad (\text{I.3.1})$$

Equivalently, the forced increment is

$$t_* = a_* - \delta = \frac{\delta b(\kappa^2 - s^2)}{\kappa(1 - b\kappa) - \delta s^2}. \quad (\text{I.3.2})$$

These bounds are sharp in the finite weighted model. If $\kappa \leq s$, no positive density increment is forced by this information alone.

Proof. The function $B_{\delta,s}(a)$ is continuous and strictly increasing for $\delta \leq a \leq 1$, after defining its endpoint value at 1 by continuity. Indeed, its nonconstant part is the secant slope $(\Phi_{\delta,s}(a) - \Phi_{\delta,s}(0))/a$. Strict convexity proves the assertion for $s > 0$; for $s = 0$ and $a \geq \delta$ it follows directly from $B_{\delta,s}(a) = 2\delta(a - \delta)/a$. Its endpoint values are

$$B_{\delta,s}(\delta) = M_\delta s, \quad B_{\delta,s}(1) = M_\delta.$$

If all $q_j \leq a$, the first chord gives $\eta \leq B_{\delta,s}(a)$. Thus the unique solution of $B_{\delta,s}(a_*) = M_\delta \kappa$ is a lower bound on $\max_j q_j$, including the endpoint $a_* = 1$ when $\kappa = 1$.

The equation for a_* is

$$\Phi_{\delta,s}(a_*) = \delta + a_*(2b\kappa - 1). \quad (\text{I.3.3})$$

Squaring, cancelling the common constant term, and removing the positive factor a_* gives a linear equation in a_* . It yields (I.3.1). There is no extraneous square-root branch. To see this explicitly, let

$$D_0 = \kappa(1 - b\kappa) - \delta s^2.$$

Since $s < \kappa$, $D_0 > \kappa(1 - b\kappa) - \delta\kappa^2 = \kappa(1 - \kappa) \geq 0$. Also

$$D_0 - \delta(\kappa - s^2) = b\kappa(1 - \kappa) \geq 0,$$

so $a_* \leq 1$, and (I.3.2) shows $a_* > \delta$. The right side of (I.3.3), after substitution, equals

$$\frac{\delta b(\kappa^2 + (1 - 2\kappa)s^2)}{D_0} > 0.$$

For $\kappa \leq 1/2$ positivity is immediate; for $\kappa > 1/2$, use $s^2 < \kappa^2$ and $1 - 2\kappa < 0$, with the case $\kappa = 1$ giving $1 - s^2 > 0$. Hence the squared solution satisfies the original root equation.

Sharpness uses densities 0 and a_* with weights $1 - \delta/a_*$ and δ/a_* and the aligned equality phases from Lemma I.1.2. Their correlation is $B_{\delta,s}(a_*) = \eta$. Finally, if $\kappa \leq s$, cells all of density δ can have correlation $M_\delta s$; opposite rotated copies reduce it to any $\eta \leq M_\delta s$ without changing any cell density. $\square$

I.3.1 Solving for the allowed arc width

At $s = 0$, the optimal increment is

$$t_0 = \frac{\delta b \kappa}{1 - b \kappa}.$$

For $0 < \kappa < 1$ and a requested increment $0 < t \leq t_0$, solving $t_*(s) = t$ gives

$$s_{\max}^2(t) = \frac{\kappa[\delta b \kappa - t(1 - b \kappa)]}{\delta(b - t)}. \quad (\text{I.3.4})$$

Here $t_0 < b$, so the denominator is positive. The right side decreases from κ^2 to 0 as t increases from 0 to t_0 . For $0 \leq s \leq s_{\max}(t)$ the forced increment is at least t . This follows either from (I.3.2), or from monotonicity of $\Phi_{\delta,s}$ in s , which makes a fixed correlation easier to obtain when the allowed arc is wider. At $\kappa = 1$, maximal correlation instead forces a density-1 cell for every $s < 1$; formula (I.3.4) need not be used at its removable endpoint.

The tradeoff has a practical interpretation. A wider allowed arc permits longer phase-flat progressions, but weakens the density increase. Formula (I.3.4) separates this exact probabilistic tradeoff from the Diophantine approximation that constructs the progressions.

I.4 Partition discrepancy and finite convex optimization

Source 01, delivered Section 5.

I.4.1 A sharp transfer using only cell discrepancy

For a partition define its normalized balanced discrepancy by

$$D = \sum_j w_j |q_j - \delta|. \quad (\text{I.4.1})$$

The equality $\sum_j w_j(q_j - \delta) = 0$ implies $D = 2\sum_j w_j(q_j - \delta)_+$, and convexity gives $D \leq M_\delta$.

Theorem I.4.1 (Discrepancy-to-density tail). *Let $0 \leq t < 1 - \delta$. Then*

$$\sum_{q_j > \delta + t} w_j \geq \max \left\{ 0, \frac{D(\delta + t) - 2\delta t}{2\delta(1 - \delta - t)} \right\}. \quad (\text{I.4.2})$$

The lower bound is sharp among all finite weighted partitions of mean δ and discrepancy D . In addition,

$$\max_j q_j \geq \delta + \frac{\delta D}{2\delta - D}. \quad (\text{I.4.3})$$

Proof. Choose a phase constant on each cell, equal to 1 or -1 according to the sign of $q_j - \delta$. Then its global correlation is exactly D . Apply Theorem I.2.1 with $s = 0$, $\eta = D$, and $a = \delta + t$. In this case

$$B_{\delta,s}(a) = \frac{2\delta t}{\delta + t}, \quad \Gamma_{\delta,s}(a) = \frac{2\delta(1 - \delta - t)}{\delta + t},$$

which gives (I.4.2). The three-density equality construction in (I.2.8) has exactly this discrepancy, proving sharpness. For (I.4.3), use Theorem I.3.1 with $s = 0$ and $\kappa = D/M_\delta$; the case $D = 0$ is just the mean bound. The denominator is positive because $D \leq 2\delta(1 - \delta) < 2\delta$. $\square$

[write] **6 October 2026; sources 11–14.** Source 13's Theorem I.27.1 is a seventh independent proof of the sharp one-sided bound, in source 02's form $R_p(t) = Q_{a,b}(p, t)$, with discarded mass and a signed drift (note there).

Suppose a uniform N -point space has a partition into m cells and a lower bound on its *average* cell size, $\bar{L} = N/m$. Let the positive right-hand side of (I.4.2) be Λ . Then some cell simultaneously has density greater than $\delta + t$ and size at least $\Lambda\bar{L}$. This is a direct local output for any supplied partition; it neither replaces an average by a minimum nor assumes that the good cell is among the largest cells in advance.

[write] **6 October 2026.** *The sharp one-sided bound, in four sources.* The bound (I.4.2) is the case $a = \delta$, $b = 1 - \delta$ of a sharp one-sided mass bound that four of the five sources prove independently, each in its own form: source 02's Theorem I.9.1 (cell means in $[-a, b]$, with the nonzero-mean extension (I.9.5)), source 03's Theorem I.14.1 (a random variable in $[-a, b]$, with the bound (I.14.1) on its essential supremum), and source 05's Lemma I.17.3 ($a = \delta$, $b = 1 - \delta$, with the non-strict event $g \geq t$, which makes its bound weaker, and a popularity corollary). The identification (checked at the write): the discrepancy is twice the positive mass, so with $T = D/2$, $a = \delta$ and $b = 1 - \delta$,

$$\frac{(a+t)T - at}{a(b-t)} = \frac{D(\delta+t) - 2\delta t}{2\delta(1-\delta-t)}, \quad \frac{aT}{a-T} = \frac{\delta D}{2\delta - D},$$

which are (I.4.2) and (I.4.3). The proof above is a genuinely different route, the case $s = 0$ of the broken-chord envelope of Theorem I.2.1; sources 02, 03 and 05 count the negative support directly. All four statements and proofs are kept. Only source 01 proves a tail bound that is sharp for a positive arc parameter $s > 0$ (Theorems I.2.1, I.2.3 and I.3.1); source 02's Theorem I.10.1, source 03's Theorem I.15.2 and source 05's Theorem I.17.4 first remove the phase and then apply the $s = 0$ bound.

[write] **6 October 2026; sources 06–10.** *Two more sources prove the same bound.* Source 06's Theorem I.21.1 states it in mass form: for cell averages in $[-a, b]$ with total mass W , positive and negative integrals $P_\pm$ and threshold η , the cells above η have mass at least $S = (P_+ - \eta(W - P_-/a))/(b - \eta)$. With $W = 1$ and $P_+ = P_- = T$ (mean zero) this is $((a + \eta)T - a\eta)/(a(b - \eta))$, source 02's bound at $t = \eta$; source 06 keeps a nonzero signed mean (its (I.21.11)) and adds the exact optimum for a fixed number of cells (note after

Corollary I.9.3). Source 10's Proposition I.23.2 is the case $a = \delta$, $b = 1 - \delta$: with $T = S_{\mathcal{P}}/2$ and $t = \tau - \delta$, the bound $((\delta + t)T - \delta t)/(\delta(1 - \delta - t))$ equals its $(\tau S_{\mathcal{P}} - 2\delta(\tau - \delta))/(2\delta(1 - \tau))$; and its (I.23.2) is the forced maximum $aT/(a - T) = \delta S_{\mathcal{P}}/(2\delta - S_{\mathcal{P}})$, i.e. (I.4.3) with $D = S_{\mathcal{P}}$ (identifications checked at the second write). Source 10's variance form (I.23.3) is not in the other sources. The bound is therefore proved independently in six sources (01, 02, 03, 05, 06, 10); all six statements and proofs are kept.

I.4.2 A master reduction for additional moment information

The preceding broken-chord argument is a small exact linear program. The same mechanism handles more density statistics without repeated pigeonholing.

Theorem I.4.2 (Finite-knot extremal reduction). *Let $0 = t_0 < t_1 < \dots < t_m = 1$, let $F : [0, 1] \rightarrow \mathbb{R}$ be convex, and let $g_1, \dots, g_r$ be continuous functions affine on each interval $[t_i, t_{i+1}]$. Fix feasible values of $\mathbb{E}g_\ell(Q)$ for a finitely supported random variable $Q \in [0, 1]$. The maximum possible $\mathbb{E}F(Q)$ is attained by a distribution supported on $\{t_0, \dots, t_m\}$. It is therefore the value of a finite linear program. An optimizer may be chosen with at most the rank of the equality-constraint matrix (including total mass) many positive weights.*

Proof. If $q \in [t_i, t_{i+1}]$, write $q = (1 - \lambda)t_i + \lambda t_{i+1}$. Replace its probability mass by masses in the proportions $1 - \lambda$ and λ at the two endpoints. Every prescribed g_ℓ -moment is unchanged, since g_ℓ is affine on this interval. The objective does not decrease, by convexity of F . Do this for every atom. The resulting optimization has finitely many nonnegative weights, one mass equation, and the given linear moment equations. Its feasible set is a closed subset of the probability simplex, so an optimum exists.

If an optimum has more positive weights than the equality matrix rank, there is a nonzero vector supported on those weights in the kernel of that matrix. Move in one of its two directions without decreasing the objective, until a weight first becomes zero. The mass equation ensures that the kernel vector has both signs, so such a move is possible. Iterate to reach the claimed support size. $\square$

Suitable moments include the mean, $|Q - \delta|$, and the stop-loss quantities $(Q - a_i)_+$, provided their breakpoints are included among the knots. The discontinuous function $\mathbf{1}_{Q > a_i}$ is not among the permitted continuous moments; its lower bound is obtained afterward from $(Q - a_i)_+/(1 - a_i) \leq \mathbf{1}_{Q > a_i}$. This distinction matters at equality.

I.4.3 Several phases on one common partition

Corollary I.4.3 (A simultaneous-phase score). *Suppose phases $\psi_1, \dots, \psi_d$ are short-arc on the same partition, with half-angles θ_i and $s_i = \sin \theta_i < 1$. For weights $\omega_i \geq 0$ with $\sum_i \omega_i = 1$, put*

$$F(q) = \sum_i \omega_i \Phi_{\delta, s_i}(q), \quad \eta = \sum_i \omega_i |\mathbb{E}(h - \delta)\psi_i|.$$

Theorem I.2.1 and its occupied-mass version hold with $\Phi_{\delta, s}$ replaced by F in the definitions of u_a , B , and Γ .

Proof. The cell bound and the triangle inequality give $\eta \leq \sum_j w_j F(q_j)$. The function F is convex, has endpoint values δ, b , and lies strictly below their chord in the interior. The proofs of Theorems I.2.1 and I.2.3 use exactly these facts. $\square$

This is an interface, not an unconditional gain over choosing the largest Fourier coefficient. Requiring one partition to control many phases has its own dimension cost. Additional relations between phases could also rule out the independent equality constructions used for the scalar weighted model.

I.5 Integer-exact simultaneous progression partitions

Source 01, delivered Section 6.

We now construct short-arc cells without asymptotic rounding conventions. An arithmetic progression here has the form $\{a, a + q, \dots, a + (\ell - 1)q\}$ with positive integer q and $\ell \geq 1$. All its terms are required to lie in $[N]_0$; no modular wrapping is used.

Theorem I.5.1 (Anisotropic phase-flat partition). *Let $\xi_1, \dots, \xi_d \in \mathbb{R}$, $d \geq 1$, and let $0 < r_i \leq 1/2$. Choose an integer $L \geq 2$ and set*

$$Q_i = \left\lceil \frac{2(L-1)}{r_i} \right\rceil, \quad T = \prod_{i=1}^d Q_i. \quad (\text{I.5.1})$$

If $N \geq LT$, there is a partition of $[N]_0$ into arithmetic progressions, all with the same positive step $q \leq T$, such that every cell has length between L and $2L - 1$ and, on every cell, each $e(\xi_i n)$ lies in an arc of length at most r_i turns.

Proof. Partition $[0, 1)^d$ into T half-open boxes with side lengths $1/Q_i$. Among the $T + 1$ points

$$(\{j\xi_1\}, \dots, \{j\xi_d\}), \quad 0 \leq j \leq T,$$

two belong to the same box. Subtracting their indices gives an integer $1 \leq q \leq T$ and integers p_i with

$$q\xi_i = p_i + \varepsilon_i, \quad |\varepsilon_i| < 1/Q_i. \quad (\text{I.5.2})$$

Partition $[N]_0$ first into residue classes modulo q . Each class has at least $\lfloor N/q \rfloor \geq L$ elements. If a class has $n = vL + w$ elements, where $v \geq 1$ and $0 \leq w < L$, divide its ordered list into $v - 1$ consecutive blocks of length L and one final block of length $L + w$. Thus every block has length in $[L, 2L - 1]$ and retains step q .

On a block of length ℓ , a real lift of the i th phase has successive differences ε_i . Its total span in turns is at most

$$(\ell - 1)|\varepsilon_i| \leq \frac{2L - 2}{Q_i} \leq r_i.$$

Projection of this real interval to the circle gives the required arc. Every point belongs to exactly one block, proving the partition assertion. $\square$

[write] **6 October 2026; sources 11–14.** Source 13's Theorem I.26.2 and Corollary I.26.3 are the same partition, proved independently; source 13 adds complete partitions with prescribed lengths $v - 1$ or v (its Theorem I.26.8) and an obstruction with independent tolerances (its Theorem I.26.5).

Corollary I.5.2 (A convenient explicit length). *For every positive integer N and the same frequencies and widths, one may take*

$$L = \max \left\{ 1, \left\lceil \left(\frac{N \prod_i r_i}{2^d} \right)^{1/(d+1)} \right\rceil \right\} \quad (\text{I.5.3})$$

and obtain cell lengths between L and $2L - 1$ with the stated arc widths. When $L = 1$, use the singleton partition with step 1.

Proof. For $L \geq 2$, the definition of Q_i gives

$$Q_i \leq \frac{2(L-1)}{r_i} + 1 \leq \frac{2L}{r_i}.$$

Consequently $L \prod_i Q_i \leq 2^d L^{d+1} / \prod_i r_i \leq N$ by (I.5.3). Apply Theorem I.5.1. Singletons satisfy the remaining case directly. $\square$

The exact condition $N \geq L \prod_i \lceil 2(L-1)/r_i \rceil$ can sometimes permit a larger L than (I.5.3). It is monotone in L and can be checked by integer arithmetic once the widths have certified rational lower bounds. The proof requires no optimal Diophantine approximation algorithm: the pigeonhole construction suffices.

[write] **6 October 2026.** Theorem I.5.1 with Corollary I.5.2 and source 05's Theorem I.18.2 are the same construction, a rectangular pigeonhole followed by residue-class blocking into lengths $[L, 2L - 1]$, with the same box counts $\lceil 2(L-1)/r_i \rceil$ and the same explicit length; credited to both. The hypotheses differ slightly: here the widths satisfy $r_i \leq 1/2$ and the theorem takes $L \geq 2$ (the corollary adds singletons), while source 05 allows widths $\varepsilon_i \leq 1$ and treats $L = 1$ inside the theorem. The one-phase case is also source 05's Proposition I.16.3 and, with a chordal error in place of an arc width, source 03's Lemma I.15.1.

I.6 A strengthened Fourier density-increment lemma

Source 01, delivered Section 7.

Theorem I.6.1 (An explicit replacement for Corollary 2.5). *Let $N \geq 1$, $h : [N]_0 \rightarrow [0, 1]$, and $\delta = N^{-1} \sum_n h(n) \in (0, 1)$. Suppose*

$$\left| \frac{1}{N} \sum_{n=0}^{N-1} (h(n) - \delta) e(\xi n) \right| \geq \eta > 0$$

for some real ξ . Put $b = 1 - \delta$ and $\kappa = \eta/(2\delta b)$. There exists a proper arithmetic progression $P \subseteq [N]_0$ such that

$$|P| \geq L := \max \left\{ 1, \left\lceil \sqrt{\frac{N}{2\pi}} \arcsin \frac{\eta}{4\delta b} \right\rceil \right\}, \quad (\text{I.6.1})$$

$$\frac{1}{|P|} \sum_{n \in P} h(n) \geq \delta + \frac{3\delta b \kappa}{4 - (4 - 3\delta)\kappa} \geq \delta + \frac{3\eta}{8}. \quad (\text{I.6.2})$$

The progression can be chosen from a partition whose cell lengths all lie between L and $2L - 1$.

Proof. As in (I.1.4), $0 < \kappa \leq 1$. Set

$$\theta = \arcsin(\kappa/2), \quad s = \sin \theta = \kappa/2, \quad r = \theta/\pi.$$

Then $0 < r \leq 1/6$. Corollary I.5.2, with one frequency, supplies the claimed partition and length. Each arc of length r turns has half-angle $\pi r = \theta$, so Theorem I.3.1 applies, also for the singleton partition. Substituting $s = \kappa/2$ into (I.3.2) gives

$$t_* = \frac{\delta b(3\kappa^2/4)}{\kappa(1 - b\kappa) - \delta\kappa^2/4} = \frac{3\delta b\kappa}{4 - (4 - 3\delta)\kappa}.$$

The denominator is positive and at most 4. Since $\eta = 2\delta b\kappa$, this is at least $3\eta/8$. $\square$

Corollary I.6.2 (Nonzero cyclic Fourier coefficients, any modulus). *Let $N \geq 2$, $A \subseteq \mathbb{Z}/N\mathbb{Z}$ have density $0 < \delta < 1$, and let $r \not\equiv 0 \pmod{N}$. Use the unnormalized transform*

$$\widehat{\mathbf{1}_A}(r) = \sum_{x \in \mathbb{Z}/N\mathbb{Z}} \mathbf{1}_A(x) e(-rx/N).$$

If $|\widehat{\mathbf{1}_A}(r)| \geq \eta N > 0$, the set of standard representatives of A has the density-increment conclusion of Theorem I.6.1 on some ambient progression $P \subseteq [N]_0$. Primality of N is unnecessary.

Proof. The geometric sum $\sum_{x=0}^{N-1} e(-rx/N)$ vanishes for any nonzero residue r , whether or not r is a unit. Thus the Fourier coefficient of $\mathbf{1}_A - \delta$ is the same as that of $\mathbf{1}_A$ at r . Apply Theorem I.6.1 with $h = \mathbf{1}_A$ and $\xi = -r/N$. $\square$

I.6.1 A simpler asymptotic comparison with all rounding exposed

Since $\arcsin u \geq u$ for $0 \leq u \leq 1$, if $N\eta \geq 32\pi\delta(1 - \delta)$ then the square root in (I.6.1) is at least 2, and

$$L \geq \sqrt{\frac{N\eta}{32\pi\delta(1 - \delta)}}. \quad (\text{I.6.3})$$

Here we used only $\lfloor x \rfloor \geq x/2$ for $x \geq 2$. Compared with the printed lower bound $L_{\text{old}} = \sqrt{\eta^3 N / (128\pi)}$, this yields

$$\frac{L}{L_{\text{old}}} \geq \frac{2}{\eta\sqrt{\delta(1 - \delta)}} \quad (\text{I.6.4})$$

in that parameter range. This is a comparison of two valid local conclusions, not a global extremal-set estimate.

δ	η	L_{old}	L from (I.6.1)	$\eta/8$	increment (I.6.2)
0.50	0.100	1576.96	126262	0.01250	0.042857
0.10	0.020	141.05	94055	0.00250	0.008359
0.01	0.002	4.46	89674	0.00025	0.000834
0.01	0.010	49.87	201570	0.00125	0.007519

Table 3: Numerical evaluations of the conditional bounds at $N = 10^{12}$. These are parameter comparisons, not constructions of sets with those Fourier coefficients. The new length is an integer; the old displayed bound is left as a real number.

[write] **6 October 2026.** *Two comparisons the source does not make. (1) The repository's proof-level baseline.* The comparison above is with the printed Corollary 2.5, which is also the statement exported by the formal project (`corollary_2_5`). The repository's proof of that corollary already obtains the stronger pair $\sqrt{\eta N / (128\pi)}$ and $\eta/4$ before weakening it (inside the proof of the private lemma `cor25_large_scale`,

Proofs02DensityIncrement.lean, line 375 at HEAD, lines 446–509; the lemma’s own statement is already the weakened pair, and it assumes $4\pi \leq \alpha N$ (corrected after the independent check of 6 October 2026; the first wording read “(the private lemma `cor25_large_scale`, ...)”); sources 03 and 05 say so (Remark I.15.3 and Section I.17, Subsection “Comparison with the actual repository baseline”). Against that baseline Theorem I.6.1 still gains: since $\arcsin u \geq u$, its length before the floor is at least $\sqrt{\eta N/(8\pi\delta(1-\delta))}$, a factor $4/\sqrt{\delta(1-\delta)}$ above $\sqrt{\eta N/(128\pi)}$, and its increment $3\eta/8$ exceeds $\eta/4$.

(2) *Cross-source comparison (intake finding, proved at the write)*. Under the same hypothesis (this η is source 05’s α and source 03’s ρ), the increment (I.6.2) is at least $3/2$ times source 05’s increment $\delta\alpha/(4\delta - \alpha)$ (Theorem I.17.4), which is also source 03’s increment at $\lambda = 1/2$ (I.15.3). *Proof*. Put $b = 1 - \delta$ and $\kappa = \eta/(2\delta b) \in (0, 1]$. Then $\delta\eta/(4\delta - \eta) = 2\delta b\kappa/(4 - 2b\kappa)$, so the ratio of the two increments is

$$\frac{3\delta b\kappa}{4 - (4 - 3\delta)\kappa} \cdot \frac{4 - 2b\kappa}{2\delta b\kappa} = \frac{3}{2} \cdot \frac{4 - 2b\kappa}{4 - (4 - 3\delta)\kappa} = \frac{3}{2} + \frac{3}{2} \cdot \frac{(2 - \delta)\kappa}{4 - (4 - 3\delta)\kappa},$$

because $(4 - 2b\kappa) - (4 - (4 - 3\delta)\kappa) = (2 - \delta)\kappa$. The denominator is at least $4 - 4\kappa + 3\delta\kappa \geq 3\delta > 0$, so the ratio exceeds $3/2$. The lengths are on the same scale: for fixed δ, η and $N \rightarrow \infty$, with $u = \eta/(4\delta b)$, the length L here is asymptotic to $\sqrt{N \arcsin(u)/(2\pi)}$, source 05’s exact L to $\sqrt{Nu/(2\pi)}$, and source 03’s length at $\lambda = 1/2$ to $\sqrt{Nu/(4\pi)}$; so this length is not asymptotically shorter than either. At finite N source 05’s exact integer length can exceed L , but by at most one: with $Y = Nu/(2\pi)$ it is $\lfloor \frac{1}{2} + \sqrt{Y + 1/4} \rfloor \leq 1 + \lfloor \sqrt{Y} \rfloor \leq L + 1$; for instance $N = 50$, $A = \{0, \dots, 24\}$, $r = 1$, $\eta = 1/(50 \sin(\pi/50)) \approx 0.3185$ give $L = 1$ here and $L = 2$ in Theorem I.17.4. Source 03’s length at $\lambda = 1/2$ never exceeds L . Among the three statements at their stated parameters, Theorem I.6.1 therefore has the largest increment, at a length of the same order. (Corrected after the independent check of 6 October 2026; the first wording read “Among the three statements at their stated parameters, Theorem I.6.1 is therefore the strongest”, which is false for lengths at finite N .) The comparison is of these stated pairs, not of trade-off curves: source 03’s parameter $\lambda \rightarrow 0$ and source 05’s width $\varepsilon \rightarrow 0$ in (I.17.9) buy increments approaching $\delta\eta/(2\delta - \eta) = \delta b\kappa/(1 - b\kappa)$, which exceeds the increment in (I.6.2) for every $\kappa < 1$ (the difference of $3(1 - b\kappa)$ and $4 - (4 - 3\delta)\kappa$ is $1 - \kappa$), at the price of shorter progressions; the loss is a constant factor near the crossing point: as $\kappa \rightarrow 0$, source 05’s trade-off reaches the increment (I.6.2) at width $\varepsilon \approx \varepsilon_0/2$, about $2^{-1/2} \approx 0.71$ of the length L here, and source 03’s at $\lambda \approx 1/4$, about $1/2$ of L (at $\kappa = 1/2$ these factors are 0.54 and 0.38; they do not depend on δ) (corrected after the independent check of 6 October 2026; the first wording read “at the price of much shorter progressions”), and source 01’s own trade-off is (I.3.4). Sources 03 and 05 are kept for these and their other features (popularity bound, exact integer length). The intake’s scan of 5,000 random (δ, α) found minimum ratio 1.500014, in agreement.

[write] 6 October 2026; sources 06–10. *Source 10’s frontier (proved at the second write)*. Source 10’s Corollary I.24.3 (its frontier at $\gamma = 1/2$, stated for sets in $\mathbb{Z}/N\mathbb{Z}$, the setting of Corollary I.6.2) has exactly the length L of (I.6.1), with increment $\delta\eta/(4\delta - \eta)$, which is source 05’s; so the increment (I.6.2) is at least $3/2$ times source 10’s at the same length (part (2) above). Source 10’s Theorem I.24.2 is a one-parameter family: for $0 < \gamma < 1$, length $L_\gamma = \max\{1, \lfloor \sqrt{N\lambda_\gamma/(2\pi)} \rfloor\}$ with $\lambda_\gamma = \arcsin((1 - \gamma)\kappa)$, and increment $D_\gamma = \delta\gamma\eta/(2\delta - \gamma\eta) = \gamma\delta b\kappa/(1 - \gamma b\kappa)$, in the notation $b = 1 - \delta$, $\kappa = \eta/(2\delta b) \in (0, 1]$ used above. *Claim*. Writing I for the increment in (I.6.2), $I \geq D_\gamma$ if and only if $\gamma(4 - \kappa) \leq 3$, and $I \geq \frac{3}{2}D_\gamma$ if and only if $\gamma(4 - (2 - \delta)\kappa) \leq 2$. *Proof*. Both denominators are positive ($4 - (4 - 3\delta)\kappa \geq 3\delta\kappa > 0$ and $1 - \gamma b\kappa > 0$), so $I \geq cD_\gamma$ is equivalent to $3(1 - \gamma b\kappa) \geq c\gamma(4 - (4 - 3\delta)\kappa)$. For $c = 1$ this is $3 - 4\gamma + \gamma\kappa(4 - 3\delta - 3b) \geq 0$, and $4 - 3\delta - 3b = 1$. For $c = 3/2$ it is $2 - 4\gamma + \gamma\kappa(4 - 3\delta - 2b) \geq 0$, and $4 - 3\delta - 2b = 2 - \delta$. Consequences: for every $\gamma \leq 1/2$ the increment here is at least $3/2$ times D_γ , while source 10’s length is at least L (because $(1 - \gamma)\kappa \geq \kappa/2$); for $1/2 < \gamma \leq 3/(4 - \kappa)$ the increment here is at least D_γ and $L \geq L_\gamma$; for $\gamma > 3/(4 - \kappa)$ source 10’s increment is larger, at a shorter length. So the intake’s “at least $3/2$ ” comparison holds for source 10 at its stated choice $\gamma = 1/2$ and for every $\gamma \leq 1/2$, not along the whole frontier. As $\gamma \rightarrow 1$, $D_\gamma \rightarrow \delta\eta/(2\delta - \eta)$, the limit described in part (2) for sources 03 and 05. (Exact rational spot checks of both equivalences, 200,000 random parameter triples, found no exception.)

[write] 6 October 2026; sources 11–14. *Source 13’s increment* (Theorem I.29.1) has exactly the length L of (I.6.1), and the increment here exceeds it by the factor $1 + 2(1 - \kappa)/(4 - (4 - 3\delta)\kappa) \in [1, 3/2]$, with equality only at $\kappa = 1$ (proved in the note after source 13’s theorem).

I.6.2 Many increment cells instead of one

On the partition in Theorem I.6.1, fix any target $\delta < a < a_*$. Then $\lambda_*(\delta, \kappa/2, a, \eta) > 0$. Theorem I.2.1 gives a disjoint family of good progressions whose union has at least $\lambda_* N$ points and whose number is at least

$$\frac{\lambda_* N}{2L - 1}.$$

Each has density greater than a and length at least L . This branching conclusion is often more useful than selecting one progression. Any further count of longer arithmetic progressions inside these cells still requires an appropriate theorem at the new density and length.

I.7 A sharp scale barrier for the phase-flat method

Source 01, delivered Section 8.

The $\sqrt{rN}$ scale of the one-frequency partition is not just an artifact of the proof. A concrete badly approximable frequency gives a matching upper scale for every individual short-arc progression.

Lemma I.7.1 (A quadratic irrational estimate). *For $\xi = (\sqrt{5} - 1)/2$ and every integer $q \geq 1$,*

$$\|q\xi\|_{\mathbb{R}/\mathbb{Z}} \geq \frac{1}{3q}. \quad (\text{I.7.1})$$

Proof. Let p be a nearest integer to $q\xi$, and set $\varphi = (1 + \sqrt{5})/2$. The nonzero integer $p^2 + pq - q^2$ is the product $(p - q\xi)(p + q\varphi)$. It is nonzero because ξ is irrational. Also $p \geq 0$ and

$$0 < p + q\varphi \leq q(\xi + \varphi) + \frac{1}{2} = q\sqrt{5} + \frac{1}{2} < 3q.$$

Taking absolute values of the product proves (I.7.1). $\square$

Theorem I.7.2 (An intrinsic square-root obstruction). *Let $0 < r < 1/2$ and $\xi = (\sqrt{5} - 1)/2$. If a proper progression $P = \{a + jq : 0 \leq j < \ell\} \subseteq [N]_0$ has all phase values $e(\xi n)$ in one arc of length r turns, then*

$$\ell \leq 1 + \sqrt{3r(N - 1)}. \quad (\text{I.7.2})$$

Proof. There is nothing to prove for $\ell = 1$. Otherwise lift every phase to the same real interval of length r . Consecutive differences of the lifts lie in $[-r, r] \subset (-1/2, 1/2)$ and are congruent to $q\xi$ modulo one. There is at most one such representative, so all these differences equal the same number ε , with $|\varepsilon| = \|q\xi\|_{\mathbb{R}/\mathbb{Z}}$. Hence

$$(\ell - 1)\|q\xi\|_{\mathbb{R}/\mathbb{Z}} \leq r.$$

Lemma I.7.1 implies $\ell - 1 \leq 3rq$. Containment in $[N]_0$ gives $q(\ell - 1) \leq N - 1$. Multiplying these two inequalities gives $(\ell - 1)^2 \leq 3r(N - 1)$. $\square$

This proves optimality of the exponent $1/2$ for a uniform theorem about single-frequency phase-flat progressions. It does not rule out longer *density-increment* progressions that are obtained by another mechanism, or longer phase-flat progressions for special frequencies.

[write] **6 October 2026.** Source 05's Theorem I.18.5 proves the same kind of obstruction for m simultaneous phases, with the algebraic vector of Lemma I.18.3; the unwrapping step is its Lemma I.18.4. At $m = 1$ its constant $c_1 = 1/20$ gives $\ell \leq 1 + \sqrt{20\varepsilon(N - 1)}$ for arcs of width ε turns, weaker than $\ell \leq 1 + \sqrt{3r(N - 1)}$ above (here $r = \varepsilon$). Credited to both; the golden-ratio constant is source 01's.

I.8 Sharp phase freezing for real functions

Source 02, delivered Section 6.

The balanced function of a set is real. This allows cancellation between its positive and negative parts to be treated more accurately than by a complex-valued triangle inequality.

Theorem I.8.1 (One-cell phase geometry). *Let P be a finite set, let $f : P \rightarrow \mathbb{R}$, and let $z : P \rightarrow \mathbb{C}$ satisfy $|z(x)| = 1$. Suppose*

$$|z(x) - z(y)| \leq 2\rho \quad (x, y \in P), \quad 0 \leq \rho \leq 1.$$

Writing $F = \sum_P f$, $L = \sum_P |f|$, and $C = |\sum_P fz|$, we have

$$C^2 \leq (1 - \rho^2)F^2 + \rho^2 L^2. \quad (\text{I.8.1})$$

Consequently

$$C \leq \rho L + (1 - \rho)|F|. \quad (\text{I.8.2})$$

The quadratic bound is sharp for any prescribed positive and negative masses and any $\rho \in [0, 1]$.

Proof. Let $p = \sum f_+$ and $q = \sum f_-$, so $F = p - q$ and $L = p + q$. If either mass is zero, the ordinary triangle inequality proves (I.8.1). Otherwise, sample a phase Z_+ with probabilities proportional to f_+ and independently sample Z_- with probabilities proportional to f_- . Then

$$\sum_P fz = \mathbb{E}(pZ_+ - qZ_-).$$

For every pair of sampled phases,

$$|pZ_+ - qZ_-|^2 = (p - q)^2 + pq|Z_+ - Z_-|^2 \leq F^2 + 4pq\rho^2 = (1 - \rho^2)F^2 + \rho^2 L^2.$$

Take expectations and use the triangle inequality. Equality is attained when all positive mass has one phase and all negative mass has another phase at distance 2ρ .

For (I.8.2), put $s = |F| \leq L$. Squaring the claimed linear majorant and subtracting the right side of (I.8.1) gives

$$[\rho L + (1 - \rho)s]^2 - [(1 - \rho^2)s^2 + \rho^2 L^2] = 2\rho(1 - \rho)s(L - s) \geq 0.$$

Both sides being nonnegative, the claim follows. $\square$

Corollary I.8.2 (Sharp aggregate freezing). *Let $P_1, \dots, P_M$ partition a finite set X , and permit a different unit-modulus phase z_j on each cell, with chordal diameter at most 2ρ . Set*

$$C = \sum_j \left| \sum_{P_j} fz_j \right|, \quad S = \sum_j \left| \sum_{P_j} f \right|, \quad L = \sum_X |f|.$$

For real f ,

$$C \leq \rho L + (1 - \rho)S. \quad (\text{I.8.3})$$

For $\rho < 1$ this gives $S \geq (C - \rho L)/(1 - \rho)$, whenever the latter lower bound is useful.

Proof. Sum (I.8.2). To see optimality of the aggregate linear law from knowledge of L, S, ρ alone, combine cells with $F = 0$ and two endpoint phases, for which $C = \rho L$, with one-sign cells having constant phase, for which $C = L = |F|$. Their contributions realize every ratio $S/L \in [0, 1]$. A zero global mean may also be imposed by using one positive and one negative one-sign cell of equal absolute mass, and rotating their phases if necessary. $\square$

For heterogeneous diameter bounds $2\rho_j$, the proof yields the more informative inequality

$$C \leq S + \sum_j \rho_j \left(\sum_{P_j} |f| - \left| \sum_{P_j} f \right| \right).$$

Thus only the part of a cell capable of positive–negative cancellation incurs a phase-freezing penalty.

Remark I.8.3 (Why the real-valued hypothesis is essential). Take phases $e^{i\theta}, e^{-i\theta}$ with $0 < \theta < \pi/2$ and take f to be their complex conjugates. Then $C = L = 2$ and $S = 2 \cos \theta$, while $\rho = \sin \theta < 1$. The right side of (I.8.3) is strictly below 2. Therefore the improved law must not be exported to arbitrary complex inputs. Gowers’s complex-disc-valued formulations still require a different estimate; the present refinement applies directly to real balanced functions.

[write] **6 October 2026.** Theorem I.8.1 is the bound printed first as source 01’s Lemma I.1.2 (note there: $\rho = s$), here for arbitrary real f and with sharpness for any prescribed masses; Corollary I.8.2 is source 01’s Lemma I.1.1. The heterogeneous inequality above and the complex-input counterexample of the Remark are source 02’s alone.

I.8.1 Balanced indicators and the correct error scale

For $f = \mathbf{1}_A - \delta$ on $|X| = N$,

$$\sum_X f = 0, \quad \frac{1}{N} \sum_X |f| = \lambda := 2\delta(1 - \delta). \quad (\text{I.8.4})$$

If $C \geq \alpha N$ and $\rho < \alpha/\lambda$, Corollary I.8.2 gives

$$\frac{S}{N} \geq \sigma(\rho) := \frac{\alpha - \rho\lambda}{1 - \rho} > 0. \quad (\text{I.8.5})$$

This improves both the use of $L \leq N$ and the loss incurred by freezing at an arbitrary phase representative. For example, at $\delta = 1/10$, $\alpha = 1/50$, and $\rho = 1/20$,

$$\lambda = \frac{9}{50}, \quad \sigma = \frac{11}{950}.$$

The crude representative-phase error bound $2\rho N = N/10$ is larger than the assumed correlation $N/50$ and is useless at these parameters, while (I.8.5) remains positive.

For a modular polynomial phase $z(x) = \exp(2\pi i \phi(x)/N)$, containment of $\phi(P)$ in a modular interval of length at most DN , with $D \leq 1/2$, gives chordal diameter at most $2 \sin(\pi D)$. The simpler estimate $\rho \leq \pi D$ will suffice in the refinement theorem.

I.9 A sharp density–mass tradeoff

Source 02, delivered Section 7.

We now isolate the selection step from all arithmetic structure.

Theorem I.9.1 (Sharp positive-tail bound). *Let $n_1, \dots, n_M > 0$ have sum N , and let $a, b > 0$. Suppose the cell means $x_j \in [-a, b]$ satisfy*

$$\sum_j n_j x_j = 0, \quad p = \frac{1}{N} \sum_j n_j \max(x_j, 0) > 0.$$

For $0 \leq t < b$, let $J_t = \{j : x_j > t\}$ and let $q_t = N^{-1} \sum_{j \in J_t} n_j$. Then

$$q_t \geq Q_{a,b}(p, t) := \frac{p(1 + t/a) - t}{b - t}. \quad (\text{I.9.1})$$

The bound is positive exactly when $t < ap/(a - p)$. Necessarily $p \leq ab/(a + b)$, so $a - p > 0$. Whenever the right side is nonnegative, the bound is attained in the weighted model by means in $\{-a, t, b\}$, allowing a zero-weight type to be omitted.

Proof. The negative mass also equals pN . Since the negative means have magnitude at most a , the total weight of negative cells is at least pN/a . Hence the total weight of positive cells is at most $N(1 - p/a)$. The cells in J_t contribute at most bq_tN positive mass, and the other positive cells contribute at most

$$t(N(1 - p/a) - q_tN).$$

Therefore $p \leq bq_t + t(1 - p/a - q_t)$, which rearranges to (I.9.1). The feasibility bound follows from $p \leq b(1 - p/a)$.

For sharpness, assign weight p/a to the mean $-a$, weight $q = Q_{a,b}(p, t)$ to b , and the remaining weight $r = 1 - p/a - q$ to t . Direct calculation gives $bq + tr = p$, so the overall mean is zero. Moreover

$$r = \frac{b(1 - p/a) - p}{b - t} \geq 0.$$

The strict threshold $x > t$ counts precisely the b -type. Thus equality holds. Rational data can be realized by integer cell sizes after scaling; arbitrary real data give a weighted sharpness statement or can be approximated by rational data. $\square$

Corollary I.9.2 (Long good cells). *Under the hypotheses of Theorem I.9.1, if $Q_{a,b}(p, t) > 0$, there is a cell with*

$$x_j > t, \quad n_j \geq \frac{Q_{a,b}(p, t)N}{M - 1}.$$

The version with M in the denominator is also valid. If every cell has size at least L_0 , there is a cell of size at least L_0 and mean at least $ap/(a - p)$.

Proof. There is at least one negative cell, and hence at most $M - 1$ good cells. Pigeonhole the good mass from (I.9.1). For the final assertion, divide the total positive mass pN by the upper bound $N(1 - p/a)$ for the total positive-cell weight. $\square$

The mass bound in Theorem I.9.1 is sharp. The particular $M - 1$ length corollary is not asserted to solve the finer extremal problem with a fixed number of cells and prescribed individual sizes.

[write] **6 October 2026.** Theorem I.9.1 is the sharp one-sided bound in its general form $[-a, b]$, proved independently in sources 02 (here), 03 (Theorem I.14.1, the same counting of the negative support in probabilistic language), 01 (the case $a = \delta$, $b = 1 - \delta$, by a convex envelope, Theorem I.4.1) and 05 (the same case with a non-strict event, Lemma I.17.3); see the note after Theorem I.4.1. The final assertion of Corollary I.9.2 (a cell of mean at least $ap/(a - p)$) is source 03's (I.14.1) and source 05's (I.17.3). The rational example of Section I.12 attains the bound.

I.9.1 A direct strengthening of Lemma 5.15

Suppose $f : X \rightarrow [-1, 1]$ has mean zero, and let x_j be its mean on a nonempty partition cell. If $S = \sum_j |\sum_{P_j} f| \geq \alpha N$ with $\alpha > 0$, then $p = S/(2N) \geq \alpha/2$ and necessarily $\alpha \leq 1$. The right side of (I.9.1) increases with p . Taking $a = b = 1$ and $t = \alpha/4$ gives the following explicit replacement.

Corollary I.9.3. *Under the hypotheses of Gowers's Lemma 5.15, after empty cells are discarded, some cell P satisfies*

$$\frac{1}{|P|} \sum_{x \in P} f(x) > \frac{\alpha}{4}, \quad |P| \geq \frac{\alpha(2 + \alpha)}{2(4 - \alpha)} \frac{N}{M - 1}. \quad (\text{I.9.2})$$

The coefficient of N/M , even without using $M - 1$, is larger than $\alpha/4$ by the factor

$$\frac{2(2 + \alpha)}{4 - \alpha} > 1.$$

For $\alpha = 2/5$ it is $2/15$, compared with the original $1/10$. The improvement follows from accounting for the negative support before spending the remaining mass on small positive means.

[write] **6 October 2026.** Source 03 states the same replacement, (I.14.4). *Comparison with the formal statement* (checked at HEAD): `lemma_5_15` (Section05.lean, line 352; exact companion `lemma_5_15_holds`, `Proofs05_10.lean`, line 735) allows $\alpha \geq 0$ and empty cells among its M cells, and concludes a cell with $\alpha|P_j|/4 \leq \sum_{P_j} f$ and $|P_j| \geq \alpha N/(4M)$. Corollary I.9.3 needs $\alpha > 0$ and counts only nonempty cells (at most M); under these conditions its mean bound is strict and its size bound is the larger one displayed. At $\alpha = 2/5$ the coefficient is $\alpha(2 + \alpha)/(2(4 - \alpha)) = 0.96/7.2 = 2/15$ (rechecked). Not formalized.

[write] **6 October 2026; sources 06–10.** *The exact optimum for a fixed number of cells (source 06).* Source 06's Theorem I.21.1 determines the largest guaranteed high cell for fixed data and a fixed number $M \geq 3$ of nonnull cells, $\min\{P_+/(b(M - 1)), S/(M - 2)\}$, with exact extremizers; its Corollary I.21.2 specializes it, under the hypotheses here ($f \in [-1, 1]$, mean zero, nonempty cells, $0 < \alpha \leq 1$), to a cell with mean $> \alpha/4$ and

$$|P_j| \geq N \min\left\{\frac{\alpha}{2(M - 1)}, \frac{S_\alpha}{M - 2}\right\}, \quad S_\alpha = \frac{\alpha(2 + \alpha)}{2(4 - \alpha)}.$$

Claim (second write). This is at least the bound $S_\alpha N/(M - 1)$ of Corollary I.9.3, strictly for $0 < \alpha < 1$, with equality at $\alpha = 1$. *Proof.* $S_\alpha/(M - 2) > S_\alpha/(M - 1)$. And $\alpha/2 \geq S_\alpha$ is equivalent to $4 - \alpha \geq 2 + \alpha$, i.e. $\alpha \leq 1$, with equality only at $\alpha = 1$; at $\alpha = 1$ the minimum is $1/(2(M - 1)) = S_1/(M - 1)$. For $M = 2$ source 06 gives $|P_j| \geq \alpha N/2 \geq S_\alpha N$. So the intake's "strictly beats" holds for $\alpha < 1$ only. Source 06's Corollary I.21.3 also improves both conclusions at once, with threshold $\tau(\alpha) = \alpha/4 + 3\alpha^2/32 + O(\alpha^3)$. Neither is formalized; `lemma_5_15` is unchanged.

I.9.2 The asymmetric range of a balanced indicator

For $f = \mathbf{1}_A - \delta$, put $a = \delta$, $b = 1 - \delta$. If $S \geq \sigma N > 0$, then every

$$0 \leq t < \frac{\delta\sigma}{2\delta - \sigma}$$

satisfies

$$\frac{1}{N} \sum_{j: |A \cap P_j|/|P_j| > \delta + t} |P_j| \geq \frac{(\sigma/2)(1 + t/\delta) - t}{1 - \delta - t}. \quad (\text{I.9.3})$$

In particular there is a cell with density at least

$$\delta + \frac{\delta\sigma}{2\delta - \sigma}. \quad (\text{I.9.4})$$

The denominator is positive because $\sigma \leq 2\delta(1 - \delta)$. The expression in (I.9.4) never exceeds 1 under this feasibility condition.

I.9.3 A nonzero-mean extension

For incomplete covers or weighted variants, mean zero can fail. If the weighted mean is μ and the weighted absolute mean is s , define $p_+ = (s + \mu)/2$ and $p_- = (s - \mu)/2$. The same proof gives

$$q_t \geq \frac{p_+ - t(1 - p_-/a)}{b - t}, \quad 0 \leq t < b. \quad (\text{I.9.5})$$

Indeed the negative support has weight at least p_-/a , while the positive mass is p_+ . This formula avoids silently using the zero-mean identity on a proper subset of the original ambient set.

I.10 Combining correlation, phase control, and density extraction

Source 02, delivered Section 8.

Theorem I.10.1 (Robust phase-to-density transfer). *Let $A \subseteq X$, $|X| = N$, $\delta = |A|/N \in (0, 1)$, and $f = \mathbf{1}_A - \delta$. Let $P_1, \dots, P_M$ be a partition into nonempty cells. Suppose each cell has a unit-modulus phase of chordal diameter at most 2ρ , where $0 \leq \rho < 1$, and*

$$\sum_j \left| \sum_{P_j} f z_j \right| \geq \alpha N, \quad \alpha > 2\rho\delta(1 - \delta).$$

Define

$$\sigma = \frac{\alpha - 2\rho\delta(1 - \delta)}{1 - \rho}, \quad Q = \frac{(\sigma/2)(1 + t/\delta) - t}{1 - \delta - t}.$$

For every $0 \leq t < \delta\sigma/(2\delta - \sigma)$, the cells with density greater than $\delta + t$ have total size at least QN , and one of them has size at least $QN/(M - 1)$. If the cells are arithmetic progressions, so is the selected cell.

Proof. Use (I.8.4) and Corollary I.8.2 to obtain $S \geq \sigma N$. Apply (I.9.3) and Corollary I.9.2. No property of a progression is used in selection, so that property is preserved automatically. $\square$

For the numerical parameters at the end of Section I.8, $\sigma = 11/950$. Formula (I.9.4) gives a density increment of at least $11/1790$. Requiring the smaller increment $t = 1/400$ guarantees good-cell mass at least $261/68210$ of the ambient set. These are exact rational consequences, not estimates obtained from a simulation.

This theorem is a replacement interface for phase-to-density arguments, rather than an unconditional restatement of every printed version of Corollary 5.8. In particular, assertions about the length of a progression must be supplied by an actual partition theorem with its thresholds and rounding conditions made explicit.

[write] **6 October 2026; sources 21–27.** Source 26’s chapter (Section I.40) supplies such a theorem at degree one, with no scale threshold: gain $\alpha/3$ at length $(N/M)^{1/3}/8$ for two-comparable parents, which implies the degree-one instance of the open `corollary_5_8`. Its degree-two sibling, source 32, and sources 41 and 42, which carry the chapter to every degree and to a square-root length, are placed and not yet written.

I.11 A uniform refinement theorem and an optimized phase budget

Source 02, delivered Section 9.

I.11.1 The imported phase-partition input

Here we make the only non-elementary input in our progression-length specialization explicit. Suppose a class of polynomial phases on injectively parameterized modular arithmetic progressions has the following property, for constants $K \geq 1$ and $R_0 \geq 1$:

Partition input $\text{PP}(K, R_0)$. For a progression of size $m > R_0$, and any integer $1 \leq v \leq m^{1/K}$, there is a partition into nonempty arithmetic subprogressions of lengths $v - 1$ or v , such that the phase values in each cell lie in a modular interval of length at most $Nm^{-1/K}$.

The supplied, edited transcription of Corollary 5.6 states this form for

$$K = (k!)^2 2^{(k+1)^2}, \quad R_0 = 2^{2^{40k^3+1}}, \quad k \geq 1, \quad (\text{I.11.1})$$

with the explicit requirement that the progression length is at most N [20]. Composing a polynomial with an injective progression chart $x \mapsto a + dx$ permits its use on each such input progression. We import this stated partition result; the lengthy Weyl-type recurrence proof is not reproved or independently Lean-verified here. The elementary theorem that follows is valid for any correctly established $\text{PP}(K, R_0)$, including stronger future inputs.

[write] **6 October 2026.** *Stale since the pin: Corollary 5.6 is now proved in Lean.* When source 02 was written (observed branch reference `66c02f8a1`), Corollary 5.6 was an open statement of the formal project, and the two sentences above were accurate. At HEAD (`25df8755d`, 6 October 2026) `corollary_5_6` (`Section05.lean`, line 205) has the exact companion `corollary_5_6_holds` (`Proofs05FullPartition.lean`, line 85). Its statement is $\text{PP}(K, R_0)$ for a polynomial ϕ on $\mathbb{Z}_N$ of degree at most k and the initial interval $\{0, \dots, r - 1\}$: $K = \text{polynomialPartitionConstant } k = (k!)^2 2^{(k+1)^2}$; the threshold $\text{polynomialPartitionThreshold } k = \text{weylThreshold } k^2 = 2^{2^{40k^3+1}} = R_0$, used with the strict inequality $R_0 < r$; $r \leq N$; $1 \leq v \leq r^{1/K}$; nonempty proper cells of length $v - 1$ or v ; image diameter at most $r^{-1/K}N$. So the Weyl-type recurrence behind (I.11.1) is now proved in the Lean development (an exact companion in the ledger). Theorem I.11.1 also composes the polynomial with the chart $x \mapsto a + dx$ of each input progression. The degree half of that step is in Lean (`polynomialOn_affine_pullback`, `Proofs05Downstream.lean`, line 244),

but Theorem I.11.1, Corollary I.11.2 and the transfer of cells back to the input progressions are not formalized. They gain no formal status. The related `corollary_5_8` is still open; `corollary_5_8_with_scale_holds` proves only its scale-qualified repair, and Section I.10 rightly declines to restate it.

Theorem I.11.1 (One global budget for many local phases). *Assume $\text{PP}(K, R_0)$. Let $Q_1, \dots, Q_M$ partition $G = \mathbb{Z}_N$ into nonempty injectively parameterized arithmetic progressions, and allow an independent polynomial phase z_i on every Q_i . Let $f : G \rightarrow \mathbb{R}$ and write*

$$C_0 = \sum_i \left| \sum_{Q_i} f z_i \right|, \quad L = \sum_G |f|.$$

For $0 < \rho < 1$, put

$$B_0 = \max\{4, (R_0 + 1)^{1/K}\}, \quad B_\rho = \max\{B_0, \pi/\rho\}. \quad (\text{I.11.2})$$

There is a refinement into J nonempty arithmetic progressions R_j such that

$$J \leq B_\rho M^{1/K} N^{1-1/K}, \quad (\text{I.11.3})$$

$$\sum_j \left| \sum_{R_j} f \right| \geq \frac{C_0 - \rho L}{1 - \rho}. \quad (\text{I.11.4})$$

Proof. Set $T_0 = \max\{4^K, R_0 + 1, (\pi/\rho)^K\} = B_\rho^K$. Treat an input cell of size m in two ways. If $m < T_0$, use singletons. Their count satisfies

$$m = m^{1-1/K} m^{1/K} \leq B_\rho m^{1-1/K},$$

and their phase diameter is zero.

If $m \geq T_0$, take $v = \lfloor m^{1/K} \rfloor$ in PP. Since $m^{1/K} \geq 4$, every resulting cell has length at least

$$v - 1 \geq m^{1/K} - 2 \geq \frac{1}{2} m^{1/K}.$$

The number of cells is at most $2m^{1-1/K}$. On each cell the half chordal diameter is at most $\pi m^{-1/K} \leq \rho$. The conditions $m > R_0$ and $m \leq N$ needed by the imported input are respected.

Summing the counts and applying concavity, or Hölder's inequality, gives

$$J \leq B_\rho \sum_i |Q_i|^{1-1/K} \leq B_\rho M^{1/K} N^{1-1/K}.$$

The sum of absolute local correlations on the refinement is at least C_0 , by the triangle inequality within each original cell. Apply Corollary I.8.2 to the refined cells and their inherited phases to obtain (I.11.4). $\square$

Unlike an argument applying a correlation threshold separately to each Q_i , the cost in (I.11.3) does not depend on a smallest positive local correlation. Cells with tiny or zero correlation cause no singular dependence. The theorem also allows different polynomials in different input cells. The exponent $1/K$ is inherited unchanged from PP.

Corollary I.11.2 (Explicit progression density increment). *In Theorem I.11.1, take $f = \mathbf{1}_A - \delta$, assume $C_0 \geq \alpha N$, and choose ρ, t as in Theorem I.10.1. With its value of Q , there is an arithmetic progression P such that*

$$\frac{|A \cap P|}{|P|} > \delta + t, \quad |P| \geq \frac{Q}{B_\rho} \left(\frac{N}{M} \right)^{1/K}. \quad (\text{I.11.5})$$

Proof. Apply Theorem I.10.1 to the refinement and use $J - 1 \leq J \leq B_\rho M^{1/K} N^{1-1/K}$. The positive lower bound for Q ensures that a negative cell and a good cell exist, so $J \geq 2$. $\square$

The word “progression” in this corollary means an injectively parameterized modular progression when the input cells are modular. Converting a wrapping modular progression to proper integer progressions is a separate cost. It must not be omitted when this interface is inserted into a proof over $[N]$.

I.11.2 Closed-form optimization

The free phase budget can be optimized exactly within the certificate just proved. Fix $\delta \in (0, 1)$, $0 < \alpha < \lambda = 2\delta(1 - \delta)$, and

$$0 \leq t < \frac{\delta\alpha}{2\delta - \alpha}.$$

Define

$$A_t = \frac{1 + t/\delta}{2}, \quad d = A_t\alpha - t, \quad e = A_t\lambda - t. \quad (\text{I.11.6})$$

The identity $e = \delta(1 - \delta - t)$ makes its positivity explicit. Then $0 < d < e$ and the good-mass certificate is

$$Q(\rho) = \frac{d - e\rho}{(1 - \rho)(1 - \delta - t)}, \quad 0 < \rho < d/e.$$

Proposition I.11.3 (Optimal uniform phase budget). *The coefficient $Q(\rho)/B_\rho$ in (I.11.5) is maximized at*

$$\boxed{\rho_{\text{opt}} = \min \left\{ 1 - \sqrt{1 - d/e}, \frac{\pi}{B_0} \right\}.} \quad (\text{I.11.7})$$

If the first entry is no larger than the second, the maximum coefficient is

$$\frac{e(1 - \sqrt{1 - d/e})^2}{\pi(1 - \delta - t)} = \frac{\delta}{\pi}(1 - \sqrt{1 - d/e})^2. \quad (\text{I.11.8})$$

Proof. When $\rho \leq \pi/B_0$, the denominator B_ρ equals π/ρ . Apart from the positive constant $1/[\pi(1 - \delta - t)]$, the objective is

$$H(\rho) = \frac{\rho(d - e\rho)}{1 - \rho}.$$

Its derivative has the sign of $d - 2e\rho + e\rho^2$. The unique critical point in $(0, d/e)$ is $r = 1 - \sqrt{1 - d/e}$; the derivative is positive before r and negative after it. At this point $d = e(2r - r^2)$, so $H(r) = er^2$.

When $\rho \geq \pi/B_0$, the denominator is constant B_0 , while $Q'(\rho) = (d - e)/[(1 - \delta - t)(1 - \rho)^2] < 0$. Hence no larger budget beyond that boundary can improve the certificate. Combining the two regions proves (I.11.7) and (I.11.8). If the boundary is outside the admissible interval, the first critical point is used automatically. $\square$

This optimization improves a particular rigorously derived lower bound; it is not a proof that the underlying polynomial partition itself is optimal. In the regime dominated by the enormous recurrence threshold R_0 , the boundary π/B_0 can control the optimum, so optimizing the elementary interface alone cannot erase the main recurrence loss.

I.12 A rational equality example for the moment bound

Source 02, delivered Appendix B.

Take $a = b = 1$, $p = 1/5$, $t = 1/10$, and assign weights

$$\frac{1}{5}, \quad \frac{2}{3}, \quad \frac{2}{15}$$

to means $-1, 1/10, 1$. Their weighted mean is

$$-\frac{1}{5} + \frac{1}{15} + \frac{2}{15} = 0,$$

the positive mass is $1/5$, and the mass of cells with mean strictly above $1/10$ is $2/15$. Formula (I.9.1) gives exactly

$$\frac{\frac{1}{5}(1 + \frac{1}{10}) - \frac{1}{10}}{1 - \frac{1}{10}} = \frac{2}{15}.$$

For $N = 150$, the three cell sizes are 30, 100, 20. Taking f constant on these cells realizes the abstract example on a finite set with integer cardinalities.

[write] 6 October 2026. The example attains Theorem I.9.1 with $a = b = 1$, $p = 1/5$, $t = 1/10$ (arithmetic rechecked at the write); it is printed here, in Part I, next to the theorem it illustrates, rather than with source 02's other appendix (Section V.11).

I.13 Conventions and analytic background

Source 03, delivered Section 2.

Unless stated otherwise, G is a finite abelian group, written additively, and $\mathbb{E}_{x \in G} = |G|^{-1} \sum_{x \in G}$. Fourier transforms in this report are normalized:

$$\widehat{f}(\chi) = \mathbb{E}_x f(x) \overline{\chi(x)}, \quad f(x) = \sum_{\chi \in \widehat{G}} \widehat{f}(\chi) \chi(x).$$

Character orthogonality says that the average of a character is zero unless it is trivial. For $f: G \rightarrow \mathbb{C}$ and $d \geq 1$, define

$$\|f\|_{U^d}^{2^d} = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{v \in \{0,1\}^d} \mathcal{C}^{|v|} f \left(x + \sum_i v_i h_i \right), \quad (\text{I.13.1})$$

where $\mathcal{C}$ denotes complex conjugation. These expressions are nonnegative. One has $\|f\|_{U^1} = |\mathbb{E}f|$, and for $d \geq 2$ they are norms, satisfy the triangle inequality, and are monotone in d . The recursive identity

$$\|f\|_{U^{d+1}}^{2^{d+1}} = \mathbb{E}_h \|f(\cdot) \overline{f(\cdot + h)}\|_{U^d}^{2^d}$$

follows by splitting the last cube coordinate. Iterated Cauchy–Schwarz proves positivity, monotonicity, and the mixed cube inequality used below; these are the standard elementary analytic facts of the source's Section 3. Fourier orthogonality also gives $\|f\|_{U^2}^4 = \sum_{\chi} |\widehat{f}(\chi)|^4$.

For $A \subset G$, write $\delta = |A|/|G|$ and $f_A = \mathbf{1}_A - \delta$. The source's statement that A is α -uniform of degree $d - 1$ means

$$\|f_A\|_{U^d}^{2^d} \leq \alpha. \quad (\text{I.13.2})$$

In particular its parameter α is a norm power, not the norm itself. This distinction is essential when comparing exponents.

A cube is a parameter tuple $(x, h_1, \dots, h_d)$ with its 2^d labelled vertices. We count parameter tuples; repeated vertices, zero increments, and coincident cubes under symmetries are retained. Arithmetic progressions in a finite group are similarly counted by their parameters (x, t) unless an ordinary integer progression is specified. A positive integer step and containment in the representative interval are proved explicitly in Section I.15.

The symbols m, N, δ are local to each section. In the cube section $m = 2^d$ is the number of vertices. In the DRC section m is the size of the witness set, while in the energy section it is $|A|$. Every statement gives its own quantifiers. An empty sum is zero and a zeroth power, including 0^0 in finite polynomial identities, is one.

I.14 An exact selection theorem for density increments

Source 03, delivered Section 3.

The last step of a density increment argument is often a finite averaging problem. A partition has appreciable total discrepancy, and one wants a cell which is both large and unusually dense. Keeping the asymmetric bounds of a balanced indicator gives more information than replacing them by $[-1, 1]$. The result below solves the associated one-sided mass problem exactly.

Theorem I.14.1 (Sharp one-sided mass bound). *Let X be a real random variable on a probability space, with $-a \leq X \leq b$, $\mathbb{E}X = 0$, where $a, b > 0$. Set*

$$T = \mathbb{E}X_+ = \mathbb{E}X_- = \frac{1}{2}\mathbb{E}|X|.$$

Then $0 \leq T \leq ab/(a+b)$. If $T > 0$, then

$$\text{ess sup } X \geq \frac{aT}{a-T}. \quad (\text{I.14.1})$$

For every $0 \leq t < b$,

$$\Pr(X > t) \geq \max \left\{ 0, \frac{(a+t)T - at}{a(b-t)} \right\}. \quad (\text{I.14.2})$$

Whenever $0 \leq t < aT/(a-T)$, the positive lower bound in (I.14.2) is attained by a probability distribution supported on $\{-a, t, b\}$. Thus it cannot be increased using only the stated information. Bound (I.14.1) is sharp as well.

Proof. Writing $p_- = \Pr(X < 0)$ gives $p_- \geq T/a$. Likewise $\Pr(X > 0) \geq T/b$, so $T/a + T/b \leq 1$. If h is the essential supremum of X , then

$$T \leq h \Pr(X > 0) \leq h(1 - T/a),$$

which proves (I.14.1). Put $q = \Pr(X > t)$. The mass where $0 < X \leq t$ is at most $1 - p_- - q$; when $t = 0$ the same inequality below still holds. Consequently

$$T \leq t(1 - p_- - q) + bq \leq t(1 - T/a) + (b - t)q.$$

Rearranging proves (I.14.2).

For sharpness set

$$p_{-a} = T/a, \quad p_b = \frac{(a+t)T - at}{a(b-t)}, \quad p_t = 1 - p_{-a} - p_b.$$

The assumed range of t makes p_b positive, and

$$p_t = \frac{ab - (a+b)T}{a(b-t)} \geq 0.$$

These weights sum to one, the negative first moment is T , and the positive first moment is $tp_t + bp_b = T$. Only the atom at b lies above t , giving equality. For (I.14.1), use the two-point distribution at $-a$ and $aT/(a-T)$, with weights T/a and $1 - T/a$. $\square$

The sharpness assertion concerns arbitrary probability weights. Rational weights give literal finite models after clearing denominators; approximation gives the corresponding limiting optimality when prescribed cardinalities or indicator-valued functions impose integrality restrictions.

[write] **6 October 2026.** Theorem I.14.1 is printed first, in cell form, as source 02's Theorem I.9.1; the proof here is the same counting of the negative support, and is kept with the statement it proves (this form adds the essential-supremum bound (I.14.1), which source 02 states as the last assertion of its Corollary I.9.2). Credited to sources 01, 02, 03 and 05 (note after source 01's Theorem I.4.1).

Corollary I.14.2 (A large cell with a specified increment). *Let $a, b > 0$, let $P_1, \dots, P_M$ be a partition of a finite probability space into cells of positive measure, and let f be real, mean zero, and valued in $[-a, b]$. Write $w_j = \mu(P_j)$ and $x_j = \mathbb{E}(f \mid P_j)$. Suppose*

$$\sum_j w_j |x_j| \geq 2T_0 > 0.$$

If $0 \leq t < aT_0/(a - T_0)$, some cell satisfies

$$x_j > t, \quad w_j \geq \frac{(a+t)T_0 - at}{a(b-t)(M-1)}. \quad (\text{I.14.3})$$

In particular $M \geq 2$ under these hypotheses.

Proof. Apply Theorem I.14.1 to the random variable taking value x_j with probability w_j . Its T is at least T_0 , and the bound in (I.14.2) is increasing in T . Since $T > 0$ and the mean is zero, at least one cell has negative average. There are therefore at most $M - 1$ cells above t . One of them carries at least $1/(M - 1)$ of their total mass. Also $T_0 \leq ab/(a + b) < a$, so the displayed denominators are positive. $\square$

For $f = 1_A - \delta$ take $a = \delta$, $b = 1 - \delta$. The resulting formula retains the density throughout the extraction. For the symmetric setting of Gowers's Lemma 5.15, put $a = b = 1$ and $2T_0 = \alpha$. At the same requested increment $t = \alpha/4$, it gives

$$\frac{|P_j|}{N} \geq \frac{\alpha(2 + \alpha)}{2(4 - \alpha)(M - 1)}, \quad \mathbb{E}(f \mid P_j) > \alpha/4, \quad (\text{I.14.4})$$

for $0 < \alpha \leq 1$. This improves the historical $\alpha/(4M)$ guarantee. More usefully, (I.14.3) allows the required increment and the retained mass to be chosen together. The exactness claim applies to the total mass formula; dividing by the number of eligible cells is a separate averaging step.

[write] **6 October 2026.** (I.14.4) is source 02's Corollary I.9.3; the comparison with the formal lemma_5_15 is in the note after it.

I.14.1 Phase removal with the actual error budget

Lemma I.14.3 (Partition phase transfer). *Let f be an integrable complex function and let P_j partition a finite probability space. On each cell let ψ_j be a phase, with $|\psi_j(x)| = 1$, and choose a constant z_j of modulus one. Put*

$$\rho = \sum_j \left| \int_{P_j} f \psi_j d\mu \right|, \quad E_{\text{osc}} = \sum_j \int_{P_j} |f| |\psi_j - z_j| d\mu.$$

Then

$$\sum_j \left| \int_{P_j} f d\mu \right| \geq \rho - E_{\text{osc}}. \quad (\text{I.14.5})$$

If f is real and mean zero, the total positive cell discrepancy is at least $(\rho - E_{\text{osc}})/2$.

Proof. For each j , subtract $z_j \int_{P_j} f$ from $\int_{P_j} f \psi_j$ and use the triangle inequality. Sum the resulting inequalities. For real f of mean zero, the positive and negative cell discrepancies have equal total mass. $\square$

This formulation allows different phases and different oscillations on different cells. Cells on which $|f|$ has small mass may tolerate a much larger phase error. If all oscillations are at most ε , then $E_{\text{osc}} \leq \varepsilon \mathbb{E}|f|$. In particular,

$$\mathbb{E}|1_A - \delta| = 2\delta(1 - \delta). \quad (\text{I.14.6})$$

For a weight $g: \Omega \rightarrow [0, 1]$ of mean δ , the same expression is an upper bound for $\mathbb{E}|g - \delta|$: pointwise, $(g - \delta)_+ \leq (1 - \delta)g$. Discarding this factor and using only $|f| \leq 1$ loses useful information, especially for small δ .

I.15 A density-sensitive Fourier increment on proper progressions

Source 03, delivered Section 4.

We give an independent progression construction, so the constants below do not depend on any unrounded version of the partition lemma in the source. All progressions in this section are ordinary progressions of integers inside $\{0, \dots, N - 1\}$, with positive integer step.

Lemma I.15.1 (Elementary phase flattening). *Let $N \geq 1$, $\theta \in \mathbb{R}$, and $0 < \varepsilon \leq 1$. Set*

$$L = \max \left\{ 1, \left\lceil \sqrt{\frac{\varepsilon N}{4\pi}} \right\rceil \right\}.$$

There is a partition of $\{0, \dots, N - 1\}$ into proper integer arithmetic progressions, each of length between L and $2L - 1$, on each of which $e(\theta x) = \exp(2\pi i \theta x)$ differs from a fixed unit complex number by at most ε .

Proof. If $L = 1$, take singletons. Otherwise set $Q = \lfloor N/L \rfloor$. The circular gaps between the $Q + 1$ points $0, \theta, \dots, Q\theta$ modulo one give an integer q with $1 \leq q \leq Q$ and

$$\|q\theta\|_{\mathbb{R}/\mathbb{Z}} \leq \frac{1}{Q + 1}.$$

This remains true if some points coincide. Split the integer interval into residue classes modulo q . Every class has at least $\lfloor N/q \rfloor \geq L$ elements. A chain of $s \geq L$ elements can be split into consecutive blocks of sizes in $[L, 2L - 1]$: write $s = hL + r$, $0 \leq r < L$, and use $h - 1$ blocks of size L and one of size $L + r$. Within any such block, the phase difference from its first point is bounded by

$$2\pi(2L - 2)\|q\theta\|_{\mathbb{R}/\mathbb{Z}} \leq \frac{4\pi(L - 1)}{Q + 1} < \frac{4\pi L^2}{N} \leq \varepsilon.$$

Here $|e(u) - e(v)| \leq 2\pi\|u - v\|_{\mathbb{R}/\mathbb{Z}}$ was used together with the triangle inequality on the circle. The blocks inherit the positive step q and remain inside the original interval. $\square$

Theorem I.15.2 (Fourier increment with density retained). *Let $A \subset \mathbb{Z}/N\mathbb{Z}$ have density $0 < \delta < 1$. Suppose a nontrivial character has normalized Fourier coefficient of magnitude at least $\rho > 0$ on 1_A . For any $0 < \lambda < 1$, there is a proper integer arithmetic progression $P \subset \{0, \dots, N - 1\}$ such that*

$$|P| \geq \max \left\{ 1, \left\lfloor \sqrt{\frac{\lambda \rho N}{8\pi\delta(1 - \delta)}} \right\rfloor \right\}, \quad (\text{I.15.1})$$

$$\frac{|A \cap P|}{|P|} \geq \delta + \frac{\delta(1 - \lambda)\rho}{2\delta - (1 - \lambda)\rho}. \quad (\text{I.15.2})$$

No primality assumption on N is needed.

Proof. Let $f = 1_A - \delta$. Orthogonality of the nontrivial character implies $\rho \leq \min\{\delta, 1 - \delta\}$ and $|\mathbb{E}_x f(x)e(-rx/N)| \geq \rho$ for the relevant r . Set

$$\varepsilon = \frac{\lambda \rho}{2\delta(1 - \delta)}.$$

Since $\max\{\delta, 1 - \delta\} \geq 1/2$, this satisfies $0 < \varepsilon \leq \lambda < 1$. Apply Lemma I.15.1 to $\theta = -r/N$. The triangle inequality over cells gives total cellwise phase correlation at least ρ . Equations (I.14.5) and (I.14.6) therefore give

$$\sum_P \mu(P) |\mathbb{E}(f | P)| \geq (1 - \lambda)\rho.$$

The positive discrepancy is at least $T_0 = (1 - \lambda)\rho/2$. Apply (I.14.1) with $a = \delta$, $b = 1 - \delta$ to obtain a cell of average at least $\delta T_0/(\delta - T_0)$. Every cell has the lower length bound (I.15.1), proving both conclusions. The denominator is positive because $\rho \leq \delta$ and $\lambda > 0$. $\square$

The balanced choice $\lambda = 1/2$ gives the particularly simple bounds

$$|P| \geq \max \left\{ 1, \left\lfloor \sqrt{\frac{\rho N}{16\pi\delta(1 - \delta)}} \right\rfloor \right\}, \quad \frac{|A \cap P|}{|P|} \geq \delta + \frac{\delta\rho}{4\delta - \rho}. \quad (\text{I.15.3})$$

The increment is at least $\rho/4$. The length keeps the factor $[\delta(1 - \delta)]^{-1/2}$, and all finite rounding is explicit. For example, if $\rho = \kappa\delta$ with $0 < \kappa \leq 1$ and $\delta \leq 1/2$, the lower length scale is $\sqrt{\kappa N/[16\pi(1 - \delta)]}$, while the density grows by the relative factor $1 + \kappa/(4 - \kappa)$. This observation is conditional on a Fourier coefficient of that relative size; it is not an assertion that every sparse set has one.

Remark I.15.3 (What was already in the repository). The pinned implementation of Corollary 2.5 already obtains an increment $\rho/4$ and length $\sqrt{\rho N/(128\pi)}$ when $\rho N \geq 4\pi$, before weakening these to its public conclusion. Thus the $\rho/4$ constant alone is not a contribution of this article. The additional conclusions here are the asymmetric selection formula, the density factor in the progression length, the adjustable loss λ , and an independent construction valid at every finite scale.

[write] **6 October 2026.** Checked at HEAD (6 October 2026): unchanged. The private `cor25_large_scale` is at `Proofs02DensityIncrement.lean`, line 375, and the exported `corollary_2_5_holds` at line 517; the file's blob (85b5c3a) is the one at both pins. *Relation to the other sources.* At $\lambda = 1/2$, (I.15.3) is the simplified form (I.17.8) of source 05's Theorem I.17.4: the same length bound and the same increment, by a different construction (a chordal error bound in place of an arc width and a disk centre). Source 05's exact integer length is asymptotically $\sqrt{2}$ times longer, and source 01's increment (Theorem I.6.1) is at least $3/2$ times this one (note after Table 3).

I.16 Exact finite partitions and the sharp rounding constant

Source 05, delivered Section 2.

I.16.1 Notation and two elementary tools

Set $I_N = \{0, 1, \dots, N-1\}$ and $\mathbb{T} = \mathbb{R}/\mathbb{Z}$. An arithmetic progression is

$$P = \{a, a+q, \dots, a+(\ell-1)q\} \subset I_N, \quad q \geq 1, \quad \ell \geq 1.$$

Its length is ℓ . All progressions considered here are proper as subsets of the integers, not just as cyclic sequences. An arc of width ε in $\mathbb{T}$ is the image of a real interval of length ε . Arc containment is a little stronger than a bound on pairwise circular distances, and it is the condition used throughout.

Lemma I.16.1 (Circular pigeonholing). *For every $\theta \in \mathbb{R}$ and integer $Q \geq 1$, there is an integer q with*

$$1 \leq q \leq Q, \quad \|q\theta\|_{\mathbb{R}/\mathbb{Z}} \leq \frac{1}{Q+1}.$$

Proof. Place $0, \theta, \dots, Q\theta$ on the circle. A repeated point gives a zero return. Otherwise, the $Q+1$ cyclic gaps have total length one, so a smallest gap has length at most $1/(Q+1)$. The difference of the two corresponding indices has absolute value between one and Q , and the circular distance of its multiple of θ is no larger than that gap. $\square$

Lemma I.16.2 (Exact residue-class blocking). *Let N, q, L be positive integers with $qL \leq N$. The interval I_N can be partitioned into progressions of common difference q , each of length between L and $2L-1$. The number of cells is at most $\lfloor N/L \rfloor$.*

Proof. Each residue class modulo q has at least L elements. If one class has $n = aL + b$ elements with $a \geq 1$ and $0 \leq b < L$, take $a-1$ consecutive blocks of length L and one final block of length $L+b$. The resulting cells have the asserted lengths. Since the cells are disjoint and each has at least L points, their number is at most N/L . $\square$

Proposition I.16.3 (Flexible one-phase partition). *Let $N \geq 1$, $0 < \varepsilon \leq 1$, and let $L \geq 1$ be an integer such that*

$$2L(L-1) \leq \varepsilon N. \tag{I.16.1}$$

Then I_N has a partition into progressions of one common positive difference, all with lengths in $[L, 2L - 1]$, on each of which $\theta x + \beta$ lies in an arc of width at most ε . This holds for every real θ, β .

Proof. Condition (I.16.1), with $\varepsilon \leq 1$, implies $L \leq N$. Put $Q = \lfloor N/L \rfloor$ and choose q using Lemma I.16.1. Then $qL \leq N$. Let $d \in [-1/2, 1/2]$ be a representative of $q\theta$ with $|d| \leq 1/(Q + 1)$. Apply Lemma I.16.2. On a cell of length $\ell \leq 2L - 1$, the phase has a real lift whose span is

$$(\ell - 1)|d| \leq \frac{2L - 2}{Q + 1} \leq \frac{2L(L - 1)}{N} \leq \varepsilon.$$

The middle inequality uses $Q + 1 > N/L$; allowing a weak inequality also covers $L = 1$. $\square$

The largest permitted integer in this proposition is

$$L_*(N, \varepsilon) = \left\lfloor \frac{1 + \sqrt{1 + 2\varepsilon N}}{2} \right\rfloor. \quad (\text{I.16.2})$$

There is no exceptional small-scale regime: $L_* = 1$ automatically covers it. With $u = \varepsilon N/2$, maximality gives $u < L_*(L_* + 1) \leq 2L_*^2$, and hence

$$L_* \geq \sqrt{\varepsilon N/4}, \quad L_* > \sqrt{\varepsilon N/2} - \frac{1}{2}. \quad (\text{I.16.3})$$

The exact integer value, rather than either weakened estimate, is preferable in implementations.

I.16.2 The optimal constant with the original upper bound

Theorem I.16.4 (Sharp capped partition). *Suppose $0 < \varepsilon \leq 1$ and $\varepsilon N \geq 1$. Every real linear phase on I_N admits a partition into arithmetic progressions, with a common positive difference, such that each image lies in an arc of width at most ε and each cell satisfies*

$$\frac{1}{3}\sqrt{\varepsilon N} \leq |P| \leq \sqrt{\varepsilon N}. \quad (\text{I.16.4})$$

The lower constant $1/3$ is the largest possible universal constant if the upper constant remains one, even if different cells are allowed different common differences.

Consequently, under the hypotheses of the corrected Gowers Lemma 2.3, its lower bound can be improved from $\sqrt{rs/(16M)}$ to

$$\sqrt{rs/(9M)}, \quad (\text{I.16.5})$$

with the same image-diameter bound s and the same upper length bound $\sqrt{rs/M}$.

Proof. Write $t = \varepsilon N$, $U = \lfloor \sqrt{t} \rfloor$, and $K = \lceil U/2 \rceil$. Then $K \geq 1$ and

$$2K(K - 1) \leq U^2 \leq t, \quad 2K - 1 \leq U.$$

Proposition I.16.3, with $L = K$, gives cells of lengths between K and $2K - 1$, and hence at most $\sqrt{t}$. Also

$$\sqrt{t} < U + 1 \leq 2K + 1 \leq 3K.$$

Thus every cell has length strictly greater than $\sqrt{t}/3$, which implies the stated weak bound.

For sharpness, take $N = 9$, a constant phase, and $\varepsilon = 1 - 1/M$ with $M \geq 10$. Then $t = 9(1 - 1/M) < 9$ and $t \rightarrow 9$ as $M \rightarrow \infty$. A putative lower constant $c > 1/3$ would eventually require every cell to have at least two points, while the upper bound $\sqrt{t} < 3$ forces every cell to have exactly two points. Nine points cannot be partitioned into pairs. This argument uses cardinalities alone and therefore applies to arbitrary allowed cell shapes.

For the modular specialization, use $N = r$, $\varepsilon = s/M$, and slope a/M for an affine map $x \mapsto ax + b \pmod{M}$. The constructed lift has integral span after multiplication by M , so its span at most s is exactly the required discrete diameter bound. The sharpness example becomes $r = 9$, $s = M - 1$, and constant affine map; it respects $r, s \leq M$ and $rs \geq M$. $\square$

Remark I.16.5. The improvement $16 \rightarrow 9$ in the denominator is an actual strengthening of the corrected statement, not a change of upper-bound convention. By contrast, using L_* in (I.16.2) gives a larger minimum but permits a larger maximum, approximately $\sqrt{2\varepsilon N}$. The two formulations serve different purposes. The sharpness of $1/3$ concerns the capped formulation only.

[write] **6 October 2026.** *Checked against the formal statement at HEAD. lemma_2_3 (Sections01_03.lean, line 141; exact companion lemma_2_3_holds, Proofs02Partition.lean, line 561) assumes $0 < N$, $0 < r$, $0 < s$, $r \leq N$, $s \leq N$, $N \leq rs$ and an affine map $x \mapsto ax + b$ from $\{0, \dots, r-1\}$ to $\mathbb{Z}_N$. It concludes a partition into proper progressions with image diameter at most s (`diameterAtMost`) and $\sqrt{rs/(16N)} \leq |P_j| \leq \sqrt{rs/N}$. Lean's N is the modulus M here. Under exactly these hypotheses $\varepsilon = s/M \leq 1$ and $\varepsilon r = rs/M \geq 1$, so Theorem I.16.4 gives the lower bound $\sqrt{rs/(9M)}$. The nine-point example ($r = 9$, $s = M - 1$, $M \geq 10$, constant map) satisfies them, since $9 \leq M$, $M - 1 \leq M$ and $M \leq 9(M - 1)$. So $1/3$ is optimal for the exact Lean formulation, as the theorem says. The improvement is not formalized.*

I.16.3 A strengthened Corollary 2.4 with the same upper length

Write $e(t) = \exp(2\pi it)$. If the values of $e(\phi(x))$ on a cell are contained in an arc of width ε , replacing the phase by its unit midpoint produces pointwise error at most $\pi\varepsilon$.

Corollary I.16.6. *Let $f : I_r \rightarrow \mathbb{C}$ satisfy $|f| \leq 1$, let $\phi(x) = \theta x + \beta$ be real linear, and suppose*

$$\left| \sum_{x \in I_r} f(x) e(-\phi(x)) \right| \geq \alpha r > 0, \quad \alpha r \geq 4\pi.$$

There is a partition into J arithmetic progressions for which

$$\sqrt{\alpha r / (36\pi)} \leq |P_j| \leq \sqrt{\alpha r / (4\pi)}, \quad (\text{I.16.6})$$

$$J \leq \sqrt{36\pi r / \alpha}, \quad (\text{I.16.7})$$

$$\sum_{j=1}^J \left| \sum_{x \in P_j} f(x) \right| \geq \frac{3}{4} \alpha r. \quad (\text{I.16.8})$$

Proof. We have $\alpha \leq 1$. Apply Theorem I.16.4 with $N = r$ and $\varepsilon = \alpha/(4\pi)$. The length and count bounds follow. On every cell choose a unit midpoint z_j for the conjugated phase arc. Then

$$\alpha r \leq \sum_j \left| \sum_{x \in P_j} f(x) \right| + \frac{\alpha}{4} \sum_x |f(x)| \leq \sum_j \left| \sum_{x \in P_j} f(x) \right| + \frac{\alpha}{4} r.$$

$\square$

Thus the corrected Corollary 2.4 constants 128π and $\alpha/2$ can be replaced by 36π and $3\alpha/4$, respectively, without changing the upper length bound. Neither a prime modulus nor a rational slope is needed. This statement does not assert optimality of the new correlation constant or of 36π .

[write] **6 October 2026.** Compared with `corollary_2_4` (`Sections01_03.lean`, line 158; exact companion `corollary_2_4_holds`, `Proofs02PhasePartition.lean`, line 276), which assumes $r \leq N$, $4\pi \leq \alpha r$ and a linear phase on $\mathbb{Z}_N$, and gives lengths in $[\sqrt{\alpha r/(128\pi)}, \sqrt{\alpha r/(4\pi)}]$, at most $\sqrt{128\pi r/\alpha}$ cells and retained correlation $(\alpha/2)r$. The corollary above keeps the upper length, improves 128π to 36π and $\alpha/2$ to $3\alpha/4$, and drops $r \leq N$, allowing any real linear phase (checked at the write against the Lean statement; not formalized).

I.17 Density-sensitive localization and popular increments

Source 05, delivered Section 3.

I.17.1 Optimal disk geometry for an arc

Lemma I.17.1 (The center need not lie on the unit circle). *Suppose a subset of the unit circle is contained in an arc of width $\varepsilon < 1/2$, with midpoint $e(\gamma)$. Set $a = \pi\varepsilon$. The disk with center $c = \cos(a)e(\gamma)$ and radius $\sin(a)$ contains the arc. No smaller disk contains both endpoints of the full arc.*

Proof. Rotate to $\gamma = 0$. For $|u| \leq a$, with $a < \pi/2$,

$$|e^{iu} - \cos a|^2 = 1 + \cos^2 a - 2 \cos a \cos u \leq \sin^2 a.$$

The endpoints have distance $2 \sin a$, giving the matching radius lower bound. $\square$

All expectations below are normalized by the total number of points. For a partition $\mathcal{P}$, let $g = \mathbb{E}(f \mid \mathcal{P})$ denote the function constant on each cell and equal to its average there.

Proposition I.17.2 (Correlation transfer with an explicit error budget). *Suppose ψ has modulus one and its values on each cell lie in an arc of width $\varepsilon < 1/2$. Let*

$$\alpha \leq |\mathbb{E}f\psi|, \quad B = \mathbb{E}|f|, \quad W = \mathbb{E}|g|.$$

Then

$$W \geq \frac{\alpha - B \sin(\pi\varepsilon)}{\cos(\pi\varepsilon)}. \quad (\text{I.17.1})$$

In particular, $W \geq \alpha - \pi\varepsilon B$ whenever this latter expression is nonnegative.

Proof. Choose the center c_P from Lemma I.17.1 on each cell. With $w_P = |P|/N$,

$$|\mathbb{E}f\psi| \leq \sum_P w_P |c_P| |\mathbb{E}_P f| + \sin(\pi\varepsilon) \mathbb{E}|f| = \cos(\pi\varepsilon) W + \sin(\pi\varepsilon) B.$$

Rearrange. The simpler bound follows from $\sin u \leq u$ and $\cos u \leq 1$. $\square$

The dependence on B is essential for sparse sets. For $f = \mathbf{1}_A - \delta$, where $\delta = |A|/N$, one has exactly

$$\mathbb{E}f = 0, \quad -\delta \leq f \leq 1 - \delta, \quad \mathbb{E}|f| = 2\delta(1 - \delta) = 2v. \quad (\text{I.17.2})$$

Bounding $\mathbb{E}|f|$ by one discards information of order δ when δ is small.

I.17.2 A sharp positive-mass conversion

Lemma I.17.3 (Capacity-sensitive extraction). *Let g be real on a finite probability space, with $\mathbb{E}g = 0$ and $-\delta \leq g \leq b = 1 - \delta$, where $0 < \delta < 1$. Put $S = \mathbb{E}g_+ = \mathbb{E}(-g)_+$. If $0 < s \leq S$, then*

$$\max g \geq \eta(s) := \frac{\delta s}{\delta - s}. \quad (\text{I.17.3})$$

For $0 < t < \eta(s)$,

$$\Pr(g \geq t) \geq \frac{s(1 + t/\delta) - t}{b - t}. \quad (\text{I.17.4})$$

In particular,

$$\Pr(g \geq \eta(s)/2) \geq \frac{s}{2(b - \eta(s)/2)}. \quad (\text{I.17.5})$$

The bound (I.17.3) is sharp as a consequence of the stated data alone.

Proof. Let $M = \max g$ and $p = \Pr(g > 0)$. Then $S \leq Mp$ and $S \leq \delta(1 - p)$. Consequently $S/M + S/\delta \leq 1$, whence $M \geq \delta S/(\delta - S)$. The right-hand side is increasing in S . Here $S \leq \delta(1 - \delta) < \delta$, so all denominators are positive.

For popularity, write $w = \Pr(g \geq t)$. The positive set has mass at most $1 - S/\delta$, and hence

$$S \leq t(1 - S/\delta) + (b - t)w.$$

Rearranging and using $S \geq s$ gives (I.17.4). Setting $t = \eta(s)/2$ gives (I.17.5). Finally, for $s \leq \delta(1 - \delta)$, the two-point distribution taking values $-\delta$ and $\eta(s)$ with probabilities s/δ and $1 - s/\delta$ has mean zero and positive mass exactly s . It attains (I.17.3). $\square$

[write] **6 October 2026.** (I.17.3) and (I.17.4) are the case $a = \delta$, $b = 1 - \delta$ of the sharp bound printed first as source 02's Theorem I.9.1, here with the event $g \geq t$ in place of $g > t$; since $\Pr(g \geq t) \geq \Pr(g > t)$, this form is weaker. Credited to sources 01, 02, 03 and 05 (note after source 01's Theorem I.4.1). The popularity bound (I.17.5) is source 05's.

I.17.3 A strengthened Fourier density-increment theorem

We normalize the Fourier transform on $\mathbb{Z}/N\mathbb{Z}$ by

$$\widehat{h}(r) = \frac{1}{N} \sum_{x=0}^{N-1} h(x)e(-rx/N).$$

The repository uses an unnormalized transform; its hypothesis $|\widehat{A}(r)| \geq \alpha N$ is our hypothesis $|\widehat{\mathbf{1}_A}(r)| \geq \alpha$.

Theorem I.17.4 (Density-sensitive Fourier increment). *Let $A \subset \mathbb{Z}/N\mathbb{Z}$ have density $0 < \delta < 1$. Suppose that for a nonzero frequency r , $|\widehat{\mathbf{1}_A}(r)| \geq \alpha > 0$. Put*

$$v = \delta(1 - \delta), \quad \varepsilon_0 = \frac{\alpha}{4\pi v}, \quad L = \left\lceil \frac{1 + \sqrt{1 + \alpha N/(2\pi v)}}{2} \right\rceil, \quad \eta = \frac{\delta \alpha}{4\delta - \alpha}. \quad (\text{I.17.6})$$

There is a partition of the standard representatives I_N into progressions of a common positive difference, all of lengths in $[L, 2L - 1]$, with the following properties:

(i) Some cell P has density at least $\delta + \eta$.

(ii) The union of cells having density at least $\delta + \eta/2$ has size at least

$$\frac{\alpha}{8(1 - \delta - \eta/2)} N. \quad (\text{I.17.7})$$

In particular, a progression exists with

$$|P| \geq \sqrt{\frac{\alpha N}{16\pi\delta(1 - \delta)}}, \quad \frac{|A \cap P|}{|P|} \geq \delta + \frac{\delta\alpha}{4\delta - \alpha} \geq \delta + \frac{\alpha}{4}. \quad (\text{I.17.8})$$

No primality assumption or lower-scale hypothesis is needed.

Proof. The sum of a nontrivial character on a finite cyclic group is zero, even when N is composite. Therefore

$$\alpha \leq \min(\delta, 1 - \delta), \quad \widehat{\mathbf{1}_A - \delta}(r) = \widehat{\mathbf{1}_A}(r).$$

In particular $\alpha/v \leq 2$, so $0 < \varepsilon_0 \leq 1/(2\pi) < 1/2$. Proposition I.16.3 gives the partition with $L = L_*(N, \varepsilon_0)$. Apply Proposition I.17.2 to $f = \mathbf{1}_A - \delta$ and the conjugated character. By (I.17.2),

$$W \geq \alpha - \pi\varepsilon_0(2v) = \alpha/2.$$

The block averages g are real and mean zero, so $\mathbb{E}g_+ = W/2 \geq \alpha/4$. Lemma I.17.3, with $s = \alpha/4$, proves (i) and (ii). Finally (I.16.3) gives $L \geq \sqrt{\varepsilon_0 N/4}$, which is exactly the first bound in (I.17.8). Since $4\delta - \alpha \leq 4\delta$, the increment is at least $\alpha/4$. $\square$

Remark I.17.5 (There is no hidden singleton exception). When the input scale is small, the formula gives $L = 1$ and the construction uses singletons. The theorem remains valid: all estimates follow from the same proof. In particular, the upper bound $\alpha \leq 1 - \delta$ prevents an impossible required density above one.

I.17.4 The full oscillation–increment trade-off

The preceding choice of ε_0 is convenient rather than optimal. For any $0 < \varepsilon < 1/2$ with $2v \sin(\pi\varepsilon) < \alpha$, the same argument gives lengths at least $L_*(N, \varepsilon)$ and positive mass at least

$$s(\varepsilon) = \frac{\alpha - 2v \sin(\pi\varepsilon)}{2 \cos(\pi\varepsilon)}. \quad (\text{I.17.9})$$

The corresponding increment is $\delta s(\varepsilon)/(\delta - s(\varepsilon))$, and the popularity formula (I.17.4) remains available. Thus length and increment can be optimized against each other at the final finite step.

For example, using

$$\varepsilon = \frac{1}{\pi} \arcsin\left(\frac{\alpha}{4v}\right)$$

gives $s(\varepsilon) = \alpha/[4\sqrt{1 - (\alpha/(4v))^2}]$, while also giving an arc width at least ε_0 . This mildly improves both quantities in Theorem I.17.4. The gain is only a constant-factor refinement; it should not be confused with the density factor $v^{-1/2}$ in the progression length.

The proof also extends to any weight $a : I_N \rightarrow [0, 1]$ of average δ , since $\mathbb{E}|a - \delta| \leq 2\delta(1 - \delta)$. The latter follows pointwise from $|u - \delta| \leq u + \delta - 2u\delta$ for $0 \leq u \leq 1$. The interpretation is then a local increase of the average weight, not of the cardinality of a subset.

I.17.5 Comparison with the actual repository baseline

For large scales, the ratio of the exact new minimum length to the stronger length already present inside the repository proof is asymptotically

$$\frac{\sqrt{\alpha N/(8\pi v)}}{\sqrt{\alpha N/(128\pi)}} = \frac{4}{\sqrt{v}}. \quad (\text{I.17.10})$$

Even the uniform square-root bound in (I.17.8) improves that baseline by a factor $\sqrt{8/v}$. These are local length comparisons under the same Fourier hypothesis.

Table 4 evaluates the formulas at $N = 10^{12}$ and $\alpha = \delta^2$. The numbers are conditional bounds, not a claim that a particular set with those parameters was constructed. Values below one in the older expression are simply vacuous real lower bounds on an integer length.

δ	Repository statement	Repository proof	New exact L	Length ratio
0.5	6233.47	24933.89	199471	8.00
0.1	49.87	4986.78	66490	13.33
0.01	0.0499	498.68	20048	40.20
0.001	0.0000499	49.87	6311	126.55

Table 4: The statement uses $\sqrt{\alpha^3 N/(128\pi)}$; the proof uses $\sqrt{\alpha N/(128\pi)}$. The last column compares the new exact integer bound with the proof-level expression.

[write] **6 October 2026.** The first row was recomputed at the write: 6233.47, 24933.89, $\lfloor (1 + \sqrt{1 + \alpha N/(2\pi v)})/2 \rfloor = 199471$ and ratio 8.00; the intake recomputed all four rows. Source 01's Theorem I.6.1 gives, under the same hypothesis, an increment at least $3/2$ times $\delta\alpha/(4\delta - \alpha)$ at a length that is not asymptotically shorter (note after Table 3); Theorem I.17.4 keeps its popularity bound (ii), its exact integer length and the trade-off (I.17.9).

I.18 Simultaneous phases and an optimal exponent

Source 05, delivered Section 4.

I.18.1 An anisotropic integer formulation

Sequentially applying a square-root localization m times can leave an exponent 2^{-m} . A simultaneous pigeonhole argument instead gives $1/(m+1)$. This is the elementary linear-phase analogue of the simultaneous localization problem that appears in more advanced density-increment arguments.

Lemma I.18.1 (Rectangular pigeonholing). *For real $\theta_1, \dots, \theta_m$ and positive integers $Q_1, \dots, Q_m$, there exists*

$$1 \leq q \leq \prod_{i=1}^m Q_i, \quad \|q\theta_i\|_{\mathbb{R}/\mathbb{Z}} < 1/Q_i \quad (1 \leq i \leq m).$$

Proof. Partition $[0, 1)^m$ into $\prod_i Q_i$ half-open rectangular boxes of side lengths $1/Q_i$. Two of the $\prod_i Q_i + 1$ points $j(\theta_1, \dots, \theta_m) \pmod{1}$ occupy the same box. Their index difference is the required q . $\square$

Theorem I.18.2 (Simultaneous linear-phase partition). *Let $N \geq 1$, $m \geq 1$, and $0 < \varepsilon_i \leq 1$. Define*

$$Q_i(L) = \max \left(1, \left\lceil \frac{2(L-1)}{\varepsilon_i} \right\rceil \right). \quad (\text{I.18.1})$$

For every positive integer L such that

$$L \prod_{i=1}^m Q_i(L) \leq N, \quad (\text{I.18.2})$$

there is a partition of I_N into progressions with one common positive difference, all of lengths in $[L, 2L-1]$, such that every phase $\theta_i x + \beta_i$ lies in an arc of width ε_i on every cell. The number of cells is at most $\lfloor N/L \rfloor$.

In particular, one may always take

$$L = \max \left(1, \left\lceil \left(\frac{N \prod_i \varepsilon_i}{2^m} \right)^{1/(m+1)} \right\rceil \right). \quad (\text{I.18.3})$$

Proof. For $L = 1$ use singletons. Otherwise apply Lemma I.18.1 with $Q_i = Q_i(L)$. The resulting return q satisfies $qL \leq N$, so Lemma I.16.2 applies. Along a cell, a real lift of phase i has span at most $(2L-2)/Q_i \leq \varepsilon_i$.

For the explicit bound, if the floor in (I.18.3) is at least two, then

$$Q_i(L) \leq \frac{2(L-1)}{\varepsilon_i} + 1 \leq \frac{2L}{\varepsilon_i},$$

and hence $L \prod_i Q_i(L) \leq 2^m L^{m+1} / \prod_i \varepsilon_i \leq N$. The cases with $L = 1$ were handled separately. $\square$

The largest integer satisfying (I.18.2) is often appreciably better than (I.18.3). Its feasibility predicate is monotone, so it can be found by binary search. The theorem has no discarded exceptional set and gives a lower bound for *every* cell, rather than merely for the average length.

[write] 6 October 2026. The same construction as source 01's Theorem I.5.1 and Corollary I.5.2 (note there); credited to both.

I.18.2 An explicit badly approximable vector

The optimality statement will be proved from scratch. This avoids importing a quantitative transference theorem with hidden constants.

Lemma I.18.3 (An algebraic obstruction). *Let $m \geq 1$, $\xi = 2^{1/(m+1)}$, and $\theta = (\xi, \xi^2, \dots, \xi^m)$. Then, for every integer $q \geq 1$,*

$$\max_{1 \leq i \leq m} \|q\xi^i\|_{\mathbb{R}/\mathbb{Z}} \geq c_m q^{-1/m}, \quad c_m = \frac{1}{2^{m+1} m (5m)^m}. \quad (\text{I.18.4})$$

Proof. Choose nearest integers p_i to $q\xi^i$ and put $h = \max_i |q\xi^i - p_i|$, $Q = \lceil q^{1/m} \rceil$. The $(Q+1)^m > q$ vectors in $\{0, \dots, Q\}^m$ cannot give distinct residues of $\sum_i u_i p_i$ modulo q . Subtract two vectors with the same residue. This gives integers a_i , not all zero, with $|a_i| \leq Q$ and $q \mid \sum_i a_i p_i$. Define

$$a_0 = -q^{-1} \sum_i a_i p_i, \quad P(X) = a_0 + \sum_{i=1}^m a_i X^i.$$

Then P is a nonzero integer polynomial of degree at most m and

$$|P(\xi)| \leq \frac{mQh}{q}. \quad (\text{I.18.5})$$

Since $\xi^i < 2$ and $|p_i/q - \xi^i| \leq 1/(2q)$, we have $|a_0| \leq (5/2)mQ$.

The polynomial $X^{m+1} - 2$ is irreducible by Eisenstein's criterion at 2. Its roots are $\xi\zeta^j$, $0 \leq j \leq m$, where $\zeta = e(1/(m+1))$. The product

$$\prod_{j=0}^m P(\xi\zeta^j)$$

is a nonzero integer: it is the algebraic norm of the algebraic integer $P(\xi)$, equivalently an integer resultant. For $j \neq 0$,

$$|P(\xi\zeta^j)| \leq |a_0| + \sum_i |a_i| \xi^i \leq 5mQ.$$

Therefore $|P(\xi)| \geq (5mQ)^{-m}$. Combining with (I.18.5) gives $h \geq q/[m(5m)^m Q^{m+1}]$. Finally $Q \leq 2q^{1/m}$ yields (I.18.4). $\square$

Lemma I.18.4 (Unwrapping a flat progression). *Suppose $0 < \varepsilon < 1/2$ and the points $\beta, \beta + \theta, \dots, \beta + (\ell - 1)\theta$ lie in an arc of width ε , with $\ell \geq 2$. Then*

$$(\ell - 1) \|\theta\|_{\mathbb{R}/\mathbb{Z}} \leq \varepsilon.$$

Proof. Choose the unique representatives of the points in a fixed real interval of length ε . Consecutive differences belong to $[-\varepsilon, \varepsilon]$ and are all congruent to θ modulo one. Since $2\varepsilon < 1$, they are the same real number. Thus the lifts form a real arithmetic progression of span $(\ell - 1) \|\theta\|_{\mathbb{R}/\mathbb{Z}}$, which is at most ε . $\square$

I.18.3 The sharp-order finite localization problem

For a vector $\theta \in \mathbb{R}^m$, let $F(\theta; N, \varepsilon)$ be the largest possible minimum cell length among all partitions of I_N into arithmetic progressions on which every $\theta_i x$ lies in an arc of width ε . Different cells may use different differences. Define

$$F_m(N, \varepsilon) = \inf_{\theta \in \mathbb{R}^m} F(\theta; N, \varepsilon).$$

This is a precise extremal formulation introduced for this note; we do not present it as a previously catalogued open problem.

Theorem I.18.5 (Optimal simultaneous localization exponent). *For $N \geq 1$ and $0 < \varepsilon < 1/2$,*

$$F_m(N, \varepsilon) \geq \max \left(1, \left\lfloor (2^{-m} \varepsilon^m N)^{1/(m+1)} \right\rfloor \right), \quad (\text{I.18.6})$$

$$F_m(N, \varepsilon) \leq 1 + (\varepsilon/c_m)^{m/(m+1)} (N - 1)^{1/(m+1)}. \quad (\text{I.18.7})$$

Consequently

$$F_m(N, \varepsilon) \asymp_m 1 + (\varepsilon^m N)^{1/(m+1)}. \quad (\text{I.18.8})$$

The exponent $1/(m+1)$ cannot be increased, even if one asks only for a single flat progression instead of a partition.

Proof. The lower bound is Theorem I.18.2 with equal widths. For the upper bound, take the explicit algebraic vector in Lemma I.18.3. If a flat progression has length $\ell \geq 2$ and difference q , then

$$q(\ell - 1) \leq N - 1, \quad c_m q^{-1/m} \leq \max_i \|q\theta_i\|_{\mathbb{R}/\mathbb{Z}} \leq \frac{\varepsilon}{\ell - 1}.$$

Rearranging gives $\ell - 1 \leq (\varepsilon/c_m)^{m/(m+1)}(N - 1)^{1/(m+1)}$. This bounds every flat progression for that vector, hence every partition minimum. Singletons satisfy the same bound. The floor and additive one in the two estimates are absorbed into constants depending only on m , proving (I.18.8). $\square$

Remark I.18.6 (Scope of sharpness). This theorem proves sharpness for linear phases and uniform arc containment. It does not prove an optimal progression length for arbitrary density increments, which might be found without flattening a prescribed phase. It also gives a symmetric-width sharpness result, not a complete anisotropic lower-bound theory for independently varying ε_i .

[write] **6 October 2026.** At $m = 1$ source 01's Theorem I.7.2 gives the sharper bound $\ell \leq 1 + \sqrt{3\varepsilon(N - 1)}$ with the golden-ratio frequency (note there).

[write] **6 October 2026; sources 11–14.** Source 13's Theorem I.26.5 treats independently varying widths $\varepsilon_i < 1/2$ by random slopes, with the matching order $(N \prod \varepsilon_i)^{1/(m+1)}$ and a finite prime-modulus certificate; loose widths remain (item I.8).

I.19 Spectral energy, residual variance, and unequal tolerances

Source 05, delivered Section 5.

I.19.1 Why several correlations should be combined

A single Fourier coefficient gives an additive increment. A collection carrying energy comparable with δ^2 can instead give an increment comparable with δ . The useful issue here is not the existence of this classical principle, but how much phase precision it needs when the input is sparse.

Let $a : I_N \rightarrow [0, 1]$ have average $\delta \in (0, 1)$, set $f = a - \delta$ and $v = \delta(1 - \delta)$. Let $r_1, \dots, r_m$ be distinct nonzero frequencies, and write

$$\chi_i(x) = e(r_i x/N), \quad c_i = \widehat{f}(r_i), \quad E = \sum_i |c_i|^2 > 0, \quad \Sigma = \sum_i |c_i|. \quad (\text{I.19.1})$$

Zero coefficients may be removed. For a partition $\mathcal{P}$, put

$$a_P = \mathbb{E}(a \mid \mathcal{P}), \quad g = a_P - \delta, \quad V = \mathbb{E}g^2.$$

Orthogonality gives $E \leq \mathbb{E}|f|^2 \leq v$, and conditional expectation gives $0 \leq V \leq v$.

Lemma I.19.1 (Residual mass remembers the captured variance). *On every cell P , $\mathbb{E}_P|a - a_P| \leq 2a_P(1 - a_P)$. Consequently*

$$\mathbb{E}|a - a_P| \leq 2(v - V). \quad (\text{I.19.2})$$

For indicator functions, equality holds.

Proof. For $0 \leq u, b \leq 1$, $|u - b| \leq u + b - 2ub$. Average with $b = a_P = \mathbb{E}_P a$. Averaging over cells gives

$$\mathbb{E} a_P (1 - a_P) = \delta - \mathbb{E} a_P^2 = \delta - (\delta^2 + V) = v - V.$$

For $u \in \{0, 1\}$ the pointwise inequality is an equality. $\square$

Theorem I.19.2 (Density-sensitive spectral extraction). *Suppose each χ_i is contained in an arc of width $\varepsilon_i < 1/2$ on every cell of $\mathcal{P}$, and set*

$$R = \sum_i |c_i| \sin(\pi \varepsilon_i). \quad (\text{I.19.3})$$

Then

$$E \leq \sqrt{EV} + 2R(v - V). \quad (\text{I.19.4})$$

If $R \leq E/(4v)$, then

$$V \geq V_0 := \frac{E}{(1 + \sqrt{1 - E/v})^2} \geq E/4, \quad (\text{I.19.5})$$

and some cell has average at least

$$\delta + V_0/\delta \geq \delta + E/(4\delta). \quad (\text{I.19.6})$$

Proof. Let $h = \sum_i c_i \chi_i$. With the normalized inner product $\langle u, v \rangle = \mathbb{E}(u\bar{v})$, orthogonality gives $\langle f, h \rangle = E$ and $\|h\|_2 = \sqrt{E}$. Decompose

$$E = \langle g, h \rangle + \langle a - a_P, h \rangle.$$

The first term has absolute value at most $\sqrt{VE}$. On each cell subtract from χ_i its disk center supplied by Lemma I.17.1. Since $a - a_P$ has zero mean on the cell, the subtraction changes no inner product. The residual term is therefore at most

$$\sum_i |c_i| \sin(\pi \varepsilon_i) \mathbb{E}|a - a_P| \leq 2R(v - V),$$

proving (I.19.4).

Assume $R \leq E/(4v)$. Since $v - V \geq 0$, divide the resulting inequality by E to obtain

$$1 \leq \sqrt{V/E} + \frac{1}{2} - \frac{V}{2v}.$$

With $x = \sqrt{V/E}$ and $\tau = E/v \in (0, 1]$, this is $\tau x^2 - 2x + 1 \leq 0$. Hence

$$x \geq \frac{1 - \sqrt{1 - \tau}}{\tau} = \frac{1}{1 + \sqrt{1 - \tau}},$$

which proves (I.19.5). Finally, if $D = \max_P a_P$, then $a_P^2 \leq D a_P$, and so $\delta^2 + V \leq D\delta$. This yields (I.19.6). $\square$

The conversion $D \geq \delta + V/\delta$ is sharp from the mean and variance alone: a random variable supported on 0 and D has variance $\delta(D - \delta)$. We do not claim that the complete phase-to-variance bound (I.19.5) is globally optimal.

I.19.2 A completely explicit progression corollary

Take equal widths

$$\varepsilon = \frac{E}{4\pi v \Sigma}. \quad (\text{I.19.7})$$

Since $E/\Sigma \leq \max_i |c_i| \leq \min(\delta, 1 - \delta)$, we have $\varepsilon \leq 1/(2\pi) < 1/2$. The inequality $\sin u \leq u$ gives $R \leq E/(4v)$. Apply Theorem I.18.2 with these widths and use the largest feasible L from (I.18.2). Then some ordinary arithmetic progression of length at least L has the increment (I.19.6). A simpler guaranteed choice is

$$L = \max \left(1, \left\lceil \left[\frac{N}{2^m} \left(\frac{E}{4\pi v \Sigma} \right)^m \right]^{1/(m+1)} \right\rceil \right). \quad (\text{I.19.8})$$

If $E \geq c\delta^2$, this gives an increment at least $c\delta/4$. Moreover $\Sigma \leq \sqrt{mE}$, so

$$\varepsilon \geq \frac{\sqrt{E}}{4\pi v \sqrt{m}} \geq \frac{\sqrt{c}}{4\pi(1-\delta)\sqrt{m}}.$$

Thus, under a fixed relative-energy hypothesis and fixed m , the required width need not shrink with δ . This is the density-sensitive feature lost by a generic unit-bound error estimate. It does not assert that every sparse progression-free set has such energy on a fixed-size spectrum.

I.19.3 Optimizing unequal phase widths

The sufficient linearized budget is

$$\sum_i b_i \varepsilon_i \leq B, \quad b_i = |c_i| > 0, \quad B = \frac{E}{4\pi v}. \quad (\text{I.19.9})$$

The guaranteed length in (I.18.3) increases with $\prod_i \varepsilon_i$. Fix a harmless cap $\varepsilon_i \leq 1/4$.

Proposition I.19.3 (Optimal allocation for the continuous bound). *If $B < \frac{1}{4} \sum_i b_i$, the product $\prod_i \varepsilon_i$ subject to (I.19.9) and $0 < \varepsilon_i \leq 1/4$ is maximized by*

$$\varepsilon_i = \min \left(\frac{1}{4}, \frac{\lambda}{b_i} \right), \quad \sum_i \min(b_i/4, \lambda) = B. \quad (\text{I.19.10})$$

If $B \geq \frac{1}{4} \sum_i b_i$, use all widths equal to $1/4$. When no cap is active,

$$\varepsilon_i = \frac{E}{4\pi m v |c_i|}, \quad \prod_i \varepsilon_i = \left(\frac{E}{4\pi m v} \right)^m \prod_i |c_i|^{-1}. \quad (\text{I.19.11})$$

Proof. Put $x_i = b_i \varepsilon_i$. Apart from a constant, the logarithm of the objective is $\sum_i \log x_i$. At a maximum with unused capacity excluded, $\sum_i x_i = B$. If two uncapped coordinates are unequal, moving a sufficiently small amount from the larger to the smaller strictly increases their product, preserves their sum, and respects both caps. Thus all uncapped x_i equal a common λ . A capped coordinate cannot exceed λ , since a small transfer from it to an uncapped coordinate would again improve the product. Consequently $x_i = \min(b_i/4, \lambda)$. Strict concavity of $\sum_i \log x_i$ makes these conditions sufficient as well; monotonicity of the displayed sum gives the unique relevant level. The uncapped formula follows immediately. $\square$

By the arithmetic–geometric mean inequality, the uncapped product improves the equal-width product by the factor

$$\frac{(\Sigma/m)^m}{\prod_i |c_i|} \geq 1.$$

This optimizes the continuous product bound, *not* the discontinuous exact ceiling problem (I.18.2). For finite implementations, evaluate the exact feasibility test after allocation; neighboring rational widths can sometimes do better. Also, adding frequencies increases energy but increases dimension. Choosing the best subset of frequencies remains a separate optimization problem.

I.19.4 An explicit extension to nonorthogonal phases

The orthogonality in Theorem I.19.2 is essential to its constants, but the argument admits a controlled extension. Let $\psi_1, \dots, \psi_m$ be arbitrary functions of modulus one on the finite probability space. Assume that $\kappa > 0$ satisfies

$$\left\| \sum_i z_i \psi_i \right\|_2^2 \leq \kappa \sum_i |z_i|^2 \quad \text{for every } (z_1, \dots, z_m) \in \mathbb{C}^m. \quad (\text{I.19.12})$$

Thus κ is a bound for the largest eigenvalue of the Gram matrix. One may always use $\kappa = m$, though a smaller certified bound is preferable. Set $c_i = \langle f, \psi_i \rangle$ and $E = \sum_i |c_i|^2 > 0$; retain v , V , and the arc-width budget R as before.

Corollary I.19.4 (Gram-controlled spectral extraction). *Under these hypotheses,*

$$E \leq \sqrt{\kappa E V} + 2R(v - V), \quad E \leq \kappa v. \quad (\text{I.19.13})$$

If $R \leq E/(4v)$, some cell has average at least

$$\delta + \frac{E}{\kappa \delta (1 + \sqrt{1 - E/(\kappa v)})^2} \geq \delta + \frac{E}{4\kappa \delta}. \quad (\text{I.19.14})$$

Proof. For $h = \sum_i c_i \psi_i$, the identity $\langle f, h \rangle = E$ still holds without orthogonality, while (I.19.12) gives $\|h\|_2 \leq \sqrt{\kappa E}$. Cauchy–Schwarz first yields $E^2 \leq \|f\|_2^2 \|h\|_2^2 \leq v \kappa E$, proving $E \leq \kappa v$. The same cellwise subtraction used in Theorem I.19.2 proves the first inequality in (I.19.13). Under the stated budget, put $x = \sqrt{\kappa V/E}$ and $\tau = E/(\kappa v)$. The resulting inequality is again $\tau x^2 - 2x + 1 \leq 0$. It gives $V \geq (E/\kappa)/(1 + \sqrt{1 - \tau})^2$. Apply $\max_P a_P \geq \delta + V/\delta$. $\square$

This includes polynomial phases and interval-restricted characters once their arc-flat partitions and a Gram bound are supplied. It does not supply either the required higher-degree partition or a useful small value of κ automatically.

I.20 Average-length partitions and exceptional-cell budgets

Source 05, delivered Section 6.

Higher-order localization statements sometimes control the number of cells rather than the minimum cell length. The following elementary interface keeps that distinction explicit and avoids unnecessarily replacing an average length by its square root.

Proposition I.20.1 (Discarding short cells quantitatively). *Let $a : I_N \rightarrow [0, 1]$ have average $\delta \in (0, 1)$, and let a partition with at most J cells satisfy*

$$\mathbb{E} |\mathbb{E}(a - \delta \mid \mathcal{P})| \geq 2s > 0.$$

For any real threshold $T > 0$ with $(1 - \delta)JT/N < s$, some cell of length at least T has average at least

$$\delta + s - (1 - \delta)JT/N. \quad (\text{I.20.1})$$

In particular, some cell has length at least

$$T = \frac{sN}{2(1 - \delta)J} \quad \text{and average at least} \quad \delta + s/2. \quad (\text{I.20.2})$$

Proof. The total positive block mass is at least s . Cells of length less than T have total normalized size at most JT/N and therefore contribute at most $(1 - \delta)JT/N$ to that positive mass. On the remaining cells the positive mass is at least the difference in (I.20.1). Since their total normalized size is at most one, one remaining cell has positive average at least that difference. $\square$

If $N/J \geq N^\gamma$, this preserves length of order sN^γ , not merely $N^{\gamma/2}$, under the specified correlation hypothesis. This is an extraction statement only. It does not justify replacing short-cell steps inside a geometric or nilpotent factorization proof, where those steps may be needed for reasons other than density extraction.

More generally, Proposition I.17.2 and Lemma I.17.3 apply to any circle-valued phase once a valid arc-flat partition is supplied. The spectral extension to an arbitrary family requires the Gram correction in Corollary I.19.4, or an orthonormal family for the uncorrected theorem. The hard issue for higher degree is supplying the appropriate partition and controlling the Gram matrix; Theorem I.18.5 cannot simply be transferred to nonlinear phases.

I.21 Sharp density extraction from a partition

Source 06, delivered Section 5.

The last step of a partition argument converts a large sum of absolute cell integrals into one cell that is both large and positively biased. Gowers' Lemma 5.15 performs this conversion by discarding cells with a small average and then cells with a small size [1]. The following result determines the exact tradeoff when the total positive and negative integrals are retained. It applies to weighted partitions, asymmetric bounds on the function, and restrictions whose mean is no longer zero.

Let a finite measure space of total mass $W > 0$ be partitioned into M nonnull cells $P_1, \dots, P_M$. Write

$$w_j = \mu(P_j), \quad x_j = \frac{1}{w_j} \int_{P_j} f d\mu, \quad -a \leq x_j \leq b,$$

where $a, b > 0$. Define the signed integral, cell variation, and its positive and negative parts by

$$m = \sum_j w_j x_j, \quad v = \sum_j w_j |x_j|, \quad P_+ = \frac{v + m}{2}, \quad P_- = \frac{v - m}{2}. \quad (\text{I.21.1})$$

The quantity v is the variation of the *cell integrals*; it need not equal $\int |f|$. Feasibility of these data implies

$$\frac{P_+}{b} + \frac{P_-}{a} \leq W. \quad (\text{I.21.2})$$

Theorem I.21.1 (Exact extraction tradeoff). *Suppose $P_+, P_- > 0$. For $0 \leq \eta < b$, set*

$$S = \frac{P_+ - \eta(W - P_-/a)}{b - \eta}. \quad (\text{I.21.3})$$

If $S > 0$, then the cells with average strictly greater than η have total mass at least S :

$$\sum_{x_j > \eta} w_j \geq S. \quad (\text{I.21.4})$$

For $M \geq 3$, at least one such cell satisfies

$$w_j \geq \min \left\{ \frac{P_+}{b(M-1)}, \frac{S}{M-2} \right\}. \quad (\text{I.21.5})$$

For $M = 2$, the corresponding bound is $w_j \geq P_+/b$.

For fixed feasible data $W, a, b, P_+, P_-, \eta, M$, the right-hand side of (I.21.5) is the exact best universal lower bound on the largest mass of a cell with $x_j > \eta$.

Proof. Let $H = \{j : x_j > \eta\}$ and $s = \sum_{j \in H} w_j$. The negative cells have total mass at least P_-/a . The positive cells outside H have average at most η , so

$$P_+ \leq bs + \eta \left(W - \frac{P_-}{a} - s \right).$$

Rearrangement proves (I.21.4). Since $P_- > 0$, there is a negative cell, and consequently $|H| \leq M-1$. If $|H| \leq M-2$, the mass bound gives a high cell of mass at least $S/(M-2)$. If $|H| = M-1$, the single remaining cell is negative. In this case all positive integral lies in H , and

$$P_+ = \sum_{j \in H} w_j x_j \leq b \sum_{j \in H} w_j.$$

One of these $M-1$ cells has mass at least $P_+/[b(M-1)]$. The same argument proves the assertion for $M = 2$.

For sharpness, it suffices to construct weighted atomic spaces. For the first branch of (I.21.5), take $M-1$ cells, each of mass $P_+/[b(M-1)]$ and average b . Give the remaining cell mass $W - P_+/b$ and average

$$-\frac{P_-}{W - P_+/b}.$$

Condition (I.21.2) places this average in $[-a, 0)$. The data are exactly the prescribed ones, and the largest high cell has the first stated mass.

For the second branch, take $M-2$ cells, each of mass $S/(M-2)$ and average b . Add a cell of average exactly η and mass

$$\ell = W - P_-/a - S,$$

and a cell of average $-a$ and mass P_-/a . The definition of S gives

$$bS + \eta\ell = P_+,$$

and (I.21.2) gives $\ell \geq 0$. If the second branch is strictly smaller than the first, then $\ell > 0$: otherwise $S = P_+/b$, which makes the second branch larger. Thus the minimizing construction has M nonnull cells and attains exactly $S/(M-2)$. At a tie the first construction suffices. This proves the asserted optimality, including the nonnull-cell requirement. $\square$

[write] **6 October 2026; sources 11–14.** This theorem answers source 13’s Question 5 (the largest guaranteed high-average cell for a fixed number of cells) in the weighted model with M fixed and the cell masses free; source 13’s Corollary I.27.2 is the weaker $M - 1$ corollary (item I.6).

The strict inequality $x_j > \eta$ is useful here: the intermediate cell in the second extremizer has average exactly η . The theorem therefore has exact extremizers, rather than only limiting ones. If only a lower bound on v is available while m is known, it can be substituted in the formulas, because both terms in (I.21.5) increase with v .

[write] **6 October 2026; sources 06–10.** *Credits.* The mass bound (I.21.4) is the sharp one-sided bound of sources 01, 02, 03, 05 and 10 in mass form (note after Theorem I.4.1); the fixed- M optimum (I.21.5) with its extremizers is proved only here. The two extremizers were rechecked at the second write: in the first, $(M - 1)$ cells at b of mass $P_+/(b(M - 1))$ carry P_+ and the last cell carries $-P_-$ on mass $W - P_+/b$; in the second, $bS + \eta\ell = P_+$ with $\ell = W - P_-/a - S \geq 0$ by (I.21.2).

I.21.1 Balanced functions and indicator functions

For a probability space, a balanced function has $W = 1$ and $m = 0$. If its cell variation is at least α , then $P_+, P_- \geq \alpha/2$, and (I.21.3) gives

$$S \geq \frac{\alpha(a + \eta) - 2a\eta}{2a(b - \eta)}. \quad (\text{I.21.6})$$

The numerator is positive precisely when

$$\eta < \frac{a\alpha}{2a - \alpha}.$$

The denominator $2a - \alpha$ is positive, since feasibility implies $\alpha \leq 2ab/(a + b) < 2a$.

Corollary I.21.2 (An explicit strengthening of Lemma 5.15). *Let f take values in $[-1, 1]$ on an N -point set with $N \geq 1$, suppose $\sum_x f(x) = 0$, and let $P_1, \dots, P_M$ be a partition into nonempty cells with*

$$\sum_{j=1}^M \left| \sum_{x \in P_j} f(x) \right| \geq \alpha N, \quad 0 < \alpha \leq 1.$$

These assumptions force $M \geq 2$. If $M \geq 3$, there is a cell such that

$$\sum_{x \in P_j} f(x) > \frac{\alpha}{4} |P_j|, \quad |P_j| \geq N \min \left\{ \frac{\alpha}{2(M - 1)}, \frac{\alpha(2 + \alpha)}{2(4 - \alpha)(M - 2)} \right\}. \quad (\text{I.21.7})$$

For $M = 2$, the size bound is $|P_j| \geq \alpha N/2$.

Proof. Use normalized counting measure, $a = b = 1$, and $\eta = \alpha/4$. Formula (I.21.6) becomes

$$S \geq S_\alpha := \frac{\alpha(2 + \alpha)}{2(4 - \alpha)}.$$

Theorem I.21.1 now gives the result. □

Each entry in the minimum in (I.21.7) is strictly larger than $\alpha/(4M)$, the size guarantee in Gowers’ lemma. There is also a simple way to improve both parameters together.

Corollary I.21.3 (Simultaneous improvement). *Under the hypotheses of Corollary I.21.2, define*

$$\tau(\alpha) = 1 - \frac{\alpha}{4} - \sqrt{1 - \alpha + \frac{\alpha^2}{16}}. \quad (\text{I.21.8})$$

Then

$$\frac{\alpha}{4} < \tau(\alpha) \leq \frac{\alpha}{2},$$

and some cell satisfies

$$\sum_{x \in P_j} f(x) > \tau(\alpha)|P_j|, \quad |P_j| \geq \frac{\tau(\alpha)}{M-1}N.$$

Moreover,

$$\tau(\alpha) = \frac{\alpha}{4} + \frac{3\alpha^2}{32} + O(\alpha^3) \quad (\alpha \rightarrow 0).$$

Proof. Put $P = \alpha/2$. The equation that makes the lower bound in (I.21.6) equal to its threshold is

$$\frac{P(1+\eta) - \eta}{1-\eta} = \eta, \quad \text{or} \quad \eta^2 + (P-2)\eta + P = 0.$$

Its smaller root is (I.21.8). The displayed interval for this root follows either by substitution at $\eta = \alpha/4$ and $\eta = \alpha/2$, or directly by squaring the nonnegative terms in (I.21.8). Thus the high cells have total mass at least $\tau(\alpha)$, and there are at most $M-1$ of them. The expansion follows by the Taylor expansion of the square root. $\square$

[write] **6 October 2026; sources 06–10.** Corollary I.21.2 is compared with the $M-1$ corollary of sources 02 and 03 in the note after Corollary I.9.3: it is at least as strong, and strictly stronger for $\alpha < 1$. Against the formal `lemma_5_15` (exact companion `lemma_5_15_holds`): that statement allows $\alpha \geq 0$ and empty cells and concludes $\alpha|P_j|/4 \leq \sum_{P_j} f$ with $|P_j| \geq \alpha N/(4M)$; the corollaries here need $\alpha > 0$ and nonempty cells. Not formalized.

For the balanced indicator $f = 1_A - \delta$ with $0 < \delta < 1$, the appropriate range is $[-\delta, 1 - \delta]$, rather than $[-1, 1]$. Keeping this asymmetry gives

$$S \geq \frac{\alpha(\delta + \eta) - 2\delta\eta}{2\delta(1 - \delta - \eta)}, \quad 0 \leq \eta < 1 - \delta. \quad (\text{I.21.9})$$

Here necessarily $\alpha \leq 2\delta(1 - \delta)$. At the maximal variation $\alpha = 2\delta(1 - \delta)$, the right-hand side of (I.21.9) equals δ for every $\eta < 1 - \delta$. This is exact: equality in the feasibility inequality forces every positive cell to have average $1 - \delta$ and every negative cell to have average $-\delta$.

The weighted sharpness examples also have an interpretation for actual subsets. For rational parameters with $a = \delta$ and $b = 1 - \delta$, replace each weighted cell by sufficiently many points and put a proportion $\delta + x_j$ of its points in A . The signed-mean condition makes the global density exactly δ . Thus the extremizers can be realized on finite uniform spaces whenever the parameters are rational; approximation gives the corresponding sharpness statement for real parameters.

I.21.2 Exceptional sets and signed drift

Restrictions of a balanced function usually cease to be balanced. The exact signed mean in (I.21.1) records this loss without inserting an additional unspecified error term.

Proposition I.21.4 (Extraction after discarding an exceptional set). *Suppose $\mu(\Omega) = 1$, $\int_{\Omega} f d\mu = 0$, and $-a \leq f \leq b$. Let the cells partition a subset of mass $1 - \rho$, and suppose their cell variation is at least α . For $0 \leq \eta \leq a$ and $\eta < b$, their high-cell mass satisfies*

$$\sum_{x_j > \eta} w_j \geq \frac{\alpha(a + \eta) - b\rho(a - \eta) - 2a\eta(1 - \rho)}{2a(b - \eta)}. \quad (\text{I.21.10})$$

More generally, a partition of mass W with $|m| \leq \varepsilon$ and $v \geq \alpha$ satisfies

$$\sum_{x_j > \eta} w_j \geq \frac{\alpha(a + \eta) - \varepsilon(a - \eta) - 2a\eta W}{2a(b - \eta)}. \quad (\text{I.21.11})$$

Proof. The argument proving (I.21.4) remains valid when one of P_+, P_- is zero. Rewriting its right-hand side in terms of v, m gives

$$S = \frac{v(a + \eta) + m(a - \eta) - 2a\eta W}{2a(b - \eta)}.$$

For the exceptional-set statement, the omitted integral belongs to $[-a\rho, b\rho]$, so its negative, the retained mean, belongs to $[-b\rho, a\rho]$. Since $a - \eta \geq 0$, substitution of $v \geq \alpha$, $m \geq -b\rho$, and $W = 1 - \rho$ proves (I.21.10). Substitution of $m \geq -\varepsilon$ proves (I.21.11). $\square$

Whenever either lower bound is positive, division by M gives a guaranteed cell size. If the observed positive and negative integrals are both nonzero, Theorem I.21.1 gives the stronger finite- M bound using their actual values. For example, $\alpha > \max(a, b)\rho$ guarantees both signs in the exceptional-set situation. These statements allow every discarded part of a structural argument to enter an explicit density budget.

[write] **6 October 2026; sources 06–10.** Exceptional sets are also treated by source 01's Theorem I.2.3 (charging only the occupied mass) and source 10's Theorem I.23.4; see the note after source 01's theorem.

I.22 Conventions and normalization

Source 10, delivered Section 2.

All group averages in this report are normalized: $\mathbb{E}_{x \in G} f(x) = |G|^{-1} \sum_{x \in G} f(x)$. For a character χ of a finite abelian group G , write

$$\widehat{f}(\chi) = \mathbb{E}_x f(x) \overline{\chi(x)}, \quad \|f\|_{L^p} = (\mathbb{E}|f|^p)^{1/p}.$$

The sum over the dual group in Parseval is an ordinary sum, so $\sum_{\chi} |\widehat{f}(\chi)|^2 = \mathbb{E}|f|^2$. If $A \subset G$ has density δ , its balanced function is $b_A = \mathbf{1}_A - \delta$. We use both indicator sets and fractional sets $a : G \rightarrow [0, 1]$, with the latter explicitly identified.

For $s \geq 1$ define the normalized Gowers cube power by

$$\|f\|_{U^s}^{2^s} = \mathbb{E}_{x, h_1, \dots, h_s} \prod_{\omega \in \{0, 1\}^s} \mathcal{C}^{|\omega|} f(x + \omega \cdot h),$$

where $\mathcal{C}$ is complex conjugation. This power is real and nonnegative: averaging one direction turns the remaining cube product into an absolute square. Thus the displayed expression defines its nonnegative 2^s -th root. In particular $\|f\|_{U^1} = \mathbb{E}|f|$ and $\|f\|_{U^2}^4 = \sum_{\chi} |\widehat{f}(\chi)|^4$.

The source uses unnormalized Fourier transforms and a degree-indexed uniformity parameter. The conversion is as follows.

Source convention	Convention in this report
$\widehat{f}_{\text{source}}(r)$	$N\widehat{f}(r)$
Unnormalized k -AP count	$N^2\Lambda_k$
u -uniform of degree d	$\ f\ _{U^{d+1}}^{2^{d+1}} \leq u$
Source gowersNorm at index $s \geq 1$	$N^{(s+1)/2^s} \ f\ _{U^s}$

Changing the sign of all difference variables changes neither cube averages nor norms. This reconciles the source's $x - h$ convention with our $x + h$ convention.

A *proper* k -term modular progression here means k distinct points $x, x + d, \dots, x + (k - 1)d$. In a general finite group, merely requiring $d \neq 0$ is insufficient. An *ordinary* progression inside $[N]_0 = \{0, \dots, N - 1\}$ is an integer progression contained in that interval; the phase-partition results produce such progressions directly. Finally, a Freiman homomorphism of order k is a map that preserves every equality between two sums of k domain elements, with repetitions allowed.

Two notational cautions prevent false gains. A claim about the average size of partition cells does not imply a lower bound for every cell, and a modular progression with distinct points is not automatically a nonwrapping integer progression. The proofs below specify the needed property each time.

I.23 Sharp density profiles on a partition

Source 10, delivered Section 3.

Let (X, μ) be a finite probability space, let $a : X \rightarrow [0, 1]$, and put $\delta = \mathbb{E}a \in (0, 1)$ and $v = \delta(1 - \delta)$. A partition $\mathcal{P}$ into cells of positive measure has cell densities $a_P = \mathbb{E}(a \mid P)$ and weights $w_P = \mu(P)$. Set

$$S_{\mathcal{P}} = \sum_{P \in \mathcal{P}} w_P |a_P - \delta|, \quad V_{\mathcal{P}} = \sum_{P \in \mathcal{P}} w_P (a_P - \delta)^2.$$

These are the first and second moments of the conditional density fluctuation. They retain information discarded by a single pigeonhole estimate.

Theorem I.23.1 (Optimal conversion of partition discrepancy). *If $M = \max_{P \in \mathcal{P}} a_P$, then*

$$S_{\mathcal{P}} \leq \frac{2\delta(M - \delta)}{M}, \tag{I.23.1}$$

$$M \geq \delta + \frac{\delta S_{\mathcal{P}}}{2\delta - S_{\mathcal{P}}}, \tag{I.23.2}$$

$$M \geq \delta + \frac{V_{\mathcal{P}}}{\delta}. \tag{I.23.3}$$

All three inequalities are sharp. Apart from zero-weight cells, equality in either moment upper bound occurs precisely when every a_P is 0 or M .

Proof. Since the mean of the a_P is δ , we have $M \geq \delta > 0$. Convexity, or a direct check on $[0, \delta]$ and $[\delta, M]$, gives

$$|u - \delta| \leq \delta + \left(1 - \frac{2\delta}{M}\right)u \quad (0 \leq u \leq M).$$

Averaging gives (I.23.1). Also $S_{\mathcal{P}} \leq 2\delta(1 - \delta) < 2\delta$, so rearrangement is legitimate and yields (I.23.2). The inequality $u^2 \leq Mu$ gives $V_{\mathcal{P}} \leq M\delta - \delta^2$, proving (I.23.3). Both pointwise inequalities are

strict in the interior of $(0, M)$, except for the harmless degenerate case $M = \delta$. In that case equality forces $u = \delta$ almost surely. A distribution giving M weight δ/M and 0 weight $1 - \delta/M$ realizes equality. For rational data this is a literal finite weighted example, and rational approximations give uniform finite examples arbitrarily close to equality. $\square$

The second-moment bound is the elementary bounded-variable variance inequality. The contribution here is its use together with the exact first-moment conversion and the geometric bounds below; neither the variance inequality nor conditional expectation is claimed as a new method. The usual estimate $M \geq \delta + S_{\mathcal{P}}/2$ loses the factor $1/(1 - S_{\mathcal{P}}/(2\delta))$.

[write] **6 October 2026; sources 06–10.** (I.23.2) is the forced maximum of the sharp one-sided bound (source 01's (I.4.3) with $D = S_{\mathcal{P}}$; source 02's and source 03's $aT/(a - T)$), and Proposition I.23.2 below is that bound itself at $a = \delta$, $b = 1 - \delta$; six sources prove it (note after Theorem I.4.1). (I.23.3) is only here.

Proposition I.23.2 (Mass of cells exceeding a density threshold). *For $\delta < \tau < 1$, let $q_{\tau} = \sum_{P:a_P > \tau} w_P$. Then*

$$q_{\tau} \geq \max \left\{ 0, \frac{\tau S_{\mathcal{P}} - 2\delta(\tau - \delta)}{2\delta(1 - \tau)} \right\}.$$

Whenever the right-hand side is positive, its dependence on $\delta, \tau, S_{\mathcal{P}}$ is sharp.

Proof. For $u \in [0, 1]$,

$$|u - \delta| \leq \delta + \left(1 - \frac{2\delta}{\tau}\right)u + \frac{2\delta(1 - \tau)}{\tau} \mathbf{1}_{\{u > \tau\}}.$$

For $u \leq \tau$ this is the chord bound on $[0, \tau]$; for $u > \tau$ the difference between the left side and the first two terms is $2\delta(u/\tau - 1)$, at most $2\delta(1 - \tau)/\tau$. Averaging and rearranging proves the claim. Equality is attained by distributions supported on $\{0, \tau, 1\}$: give 1 weight q , τ weight $(\delta - q)/\tau$, and 0 the remaining weight. As $0 \leq q \leq \delta$ varies, these realize the entire positive-bound range. $\square$

I.23.1 Phase approximation with a density-sensitive error budget

Write $f = a - \delta$. Its mean is zero, and

$$\mathbb{E}|f| \leq 2v, \quad \mathbb{E}|f|^2 \leq v. \quad (\text{I.23.4})$$

Both are equalities for an indicator $a = \mathbf{1}_A$. The first follows from the chord inequality $|u - \delta| \leq \delta + (1 - 2\delta)u$ on $[0, 1]$; the second follows from $u^2 \leq u$.

Theorem I.23.3 (Correlation-to-density transfer). *Let $\phi : X \rightarrow \mathbb{C}$ be 1-bounded, and suppose $|\mathbb{E}f\phi| \geq \eta > 0$. Suppose that for each $P \in \mathcal{P}$ there is z_P in the closed unit disc such that $|\phi(x) - z_P| \leq \rho$ for $x \in P$. If $s = \eta - 2v\rho > 0$, then some cell satisfies*

$$a_P \geq \delta + \frac{\delta s}{2\delta - s}.$$

The tail bound of Proposition I.23.2 also holds with $S_{\mathcal{P}}$ replaced by s .

Proof. The triangle inequality and (I.23.4) give

$$\eta \leq \left| \sum_P z_P \int_P f d\mu \right| + \rho \mathbb{E}|f| \leq S_{\mathcal{P}} + 2v\rho.$$

Apply Theorem I.23.1 and Proposition I.23.2. $\square$

This is valid for arbitrary bounded phases, arbitrary finite weights, and fractional sets $a \in [0, 1]$. Its improvement over a crude 1-bounded argument is substantial at small density: a phase error of size ρ costs at most $2\delta(1 - \delta)\rho$, rather than ρ .

Theorem I.23.4 (Allowance for an exceptional set). *In Theorem I.23.3, suppose instead that the cells partition $X \setminus E$, where $\mu(E) = \varepsilon < \delta$, and that the phase approximation holds only there. Put $s = \eta - 2v\rho$ and assume $s > 2(1 - \delta)\varepsilon$. Then a retained cell satisfies*

$$a_P \geq \delta + \frac{\delta(s - 2(1 - \delta)\varepsilon)}{2\delta(1 - \varepsilon) - s}. \quad (\text{I.23.5})$$

Proof. Let $b = \int_E a \, d\mu \leq \varepsilon$. The chord inequality gives $\int_E |f| \, d\mu \leq \delta\varepsilon + (1 - 2\delta)b$. Writing S for the unnormalized discrepancy over retained cells, the correlation estimate gives $s \leq S + \int_E |f| \, d\mu$. If all retained densities were at most δ , the right side would be at most $2(1 - \delta)b$, contradicting the hypothesis. Thus their maximum M exceeds δ . The chord bound on $[0, M]$, now with retained mass $1 - \varepsilon$ and retained a -mass $\delta - b$, gives

$$s \leq 2\delta(1 - b) - \frac{2\delta(\delta - b)}{M}.$$

Consequently

$$M \geq \frac{2\delta(\delta - b)}{2\delta(1 - b) - s} \geq \frac{2\delta(\delta - \varepsilon)}{2\delta(1 - \varepsilon) - s}.$$

The first fraction is nonincreasing in b , because its derivative has the sign of $s - 2v \leq 0$. All denominators are positive: $s \leq 2v$ and $\varepsilon < \delta$. Subtracting δ proves the formula. $\square$

The exceptional-set formula makes the usable error budget explicit. An exceptional proportion ε can consume all guaranteed positive imbalance unless the residual correlation exceeds $2(1 - \delta)\varepsilon$. The formula concerns the transfer step; it does not assert optimality of a preceding geometric construction.

[write] **6 October 2026; sources 06–10.** Compare source 01's Theorem I.2.3, which charges only the occupied mass of the exceptional set, and source 06's Proposition I.21.4; the three are different models (note after source 01's theorem).

I.24 A constructive Fourier-to-progression tradeoff

Source 10, delivered Section 4.

Write $[N]_0 = \{0, \dots, N - 1\}$ and $e(t) = \exp(2\pi it)$. The next partition lemma avoids a large-scale side condition and keeps all integer rounding visible.

Lemma I.24.1 (An explicit phase partition). *Let $N \geq 1$, $\theta \in \mathbb{R}$, and $0 < \lambda \leq \pi/2$. Define*

$$L = \max \left\{ 1, \left\lceil \sqrt{\frac{N\lambda}{2\pi}} \right\rceil \right\}.$$

There is a partition of $[N]_0$ into ordinary integer arithmetic progressions, each of length between L and $2L - 1$, such that on every cell P the phase $e(\theta x)$ lies within distance $\sin \lambda$ of a complex number z_P with $|z_P| \leq 1$.

Proof. For $L = 1$, take singletons. Otherwise set $Q = \lfloor N/L \rfloor$. Among the $Q+1$ points $0, \theta, \dots, Q\theta$ on $\mathbb{R}/\mathbb{Z}$, two have circular distance at most $1/(Q+1)$: sort them around the circle and use the smallest circular gap, treating repeated points separately. Thus there is $1 \leq q \leq Q$ with $\|q\theta\|_{\mathbb{R}/\mathbb{Z}} \leq 1/(Q+1)$. Each residue class modulo q inside $[N]_0$ has at least $\lfloor N/q \rfloor \geq L$ points. A string of $n \geq L$ points can be cut into blocks of sizes between L and $2L-1$: write $n = aL + b$ with $a \geq 1$ and $0 \leq b < L$, use one block of length $L+b$, and $a-1$ blocks of length L .

Choose a signed representative t of $q\theta$ with $|t| \leq 1/(Q+1)$. On a block of $m \leq 2L-1$ consecutive points of a residue class, the phase lies on an arc centered at its midpoint phase u_P , with half-angle

$$H = \pi(m-1)|t| \leq \frac{2\pi(L-1)}{Q+1} < \frac{2\pi L^2}{N} \leq \lambda.$$

Set $z_P = (\cos H)u_P$. For a point on the arc making angle h with u_P , where $|h| \leq H \leq \pi/2$,

$$|e^{ih} - \cos H|^2 = 1 + \cos^2 H - 2 \cos H \cos h \leq \sin^2 H.$$

This proves the claimed disc approximation and all the length bounds. $\square$

Theorem I.24.2 (A one-parameter density-length frontier). *Let $a : [N]_0 \rightarrow [0, 1]$ have mean $\delta \in (0, 1)$, put $v = \delta(1 - \delta)$, and suppose*

$$\left| \frac{1}{N} \sum_{x=0}^{N-1} (a(x) - \delta)e(\theta x) \right| \geq \eta > 0.$$

For every $0 < \gamma < 1$, set

$$\begin{aligned} \lambda_\gamma &= \arcsin\left(\frac{(1-\gamma)\eta}{2v}\right), \\ L_\gamma &= \max\left\{1, \left\lfloor \sqrt{\frac{N\lambda_\gamma}{2\pi}} \right\rfloor\right\}, \\ D_\gamma &= \frac{\delta\gamma\eta}{2\delta - \gamma\eta}. \end{aligned}$$

There exists an ordinary arithmetic progression $P \subset [N]_0$ with

$$L_\gamma \leq |P| \leq 2L_\gamma - 1, \quad \frac{1}{|P|} \sum_{x \in P} a(x) \geq \delta + D_\gamma.$$

Moreover, the cells of the phase partition whose density exceeds any $\tau \in (\delta, \delta + D_\gamma)$ cover at least

$$N \frac{\tau\gamma\eta - 2\delta(\tau - \delta)}{2\delta(1 - \tau)}$$

points. Their number is at least this quantity divided by $2L_\gamma - 1$.

Proof. By (I.23.4), $\eta \leq 2v$, so the arcsine is defined, strictly positive, and less than $\pi/2$. Apply Lemma I.24.1 with $\lambda = \lambda_\gamma$ and Theorem I.23.3 with $\rho = \sin \lambda_\gamma = (1 - \gamma)\eta/(2v)$. Then $s = \gamma\eta$. Proposition I.23.2 gives the final assertion. $\square$

Corollary I.24.3 (A convenient balanced choice). *If $A \subset \mathbb{Z}/N\mathbb{Z}$ has density $\delta \in (0, 1)$ and some nontrivial normalized Fourier coefficient has modulus at least $\eta > 0$, then A , represented inside $[N]_0$, has a progression with density at least*

$$\delta + \frac{\delta\eta}{4\delta - \eta}$$

and length at least

$$\max \left\{ 1, \left\lfloor \sqrt{\frac{N}{2\pi} \arcsin\left(\frac{\eta}{4\delta(1-\delta)}\right)} \right\rfloor \right\}. \quad (\text{I.24.1})$$

No primality assumption on N is required.

The corresponding intermediate estimate already in the repository’s proof of Corollary 2.5 is length at least $(\eta N/(128\pi))^{1/2}$ and increment $\eta/4$, in its stated large-scale range. Thus merely replacing the printed η^3 by η would not be a new improvement relative to the supplied version. Our additional gain is the explicit density factor, the optimized cap conversion, improved geometric constants, and the unrestricted finite-scale statement. For $\eta/v \rightarrow 0$, the continuous expression underlying (I.24.1) is

$$\sqrt{\frac{\eta N}{8\pi v}} \left(1 + O((\eta/v)^2) \right).$$

Its ratio to the repository’s continuous lower bound is $(4/\sqrt{v})(1 + O((\eta/v)^2))$; floors may matter at small N . The density gain is $\eta/[4(1 - \eta/(4\delta))] \geq \eta/4$.

I.24.1 Numerical scale of the guarantees

Figure I.24.1 evaluates the explicit family in Theorem I.24.2. Increasing the retained correlation fraction γ increases the guaranteed density increment and reduces the minimum length supplied by this partition construction. The curve is a family of sufficient guarantees; optimality of its geometric constants is not asserted.

For a concrete comparison, the following table uses $\gamma = 1/2$ and $N = 10^8$ throughout. The older length column evaluates the continuous lower bound already present in the supplied proof; the new length column includes the floor in (I.24.1). All three parameter choices satisfy the older large-scale condition.

δ	η	Older length	New length	Older gain	New gain
0.5	0.05	111.51	892	0.0125	0.0128205
0.1	0.01	49.87	664	0.0025	0.00256410
0.01	0.0001	4.99	200	0.000025	0.0000250627

The improvement in the length guarantee is especially visible at low density because the phase approximation error is charged to $\mathbb{E}|\mathbf{1}_A - \delta| = 2\delta(1 - \delta)$. Each entry is conditional on the presence of the specified Fourier coefficient; the table does not claim that an arbitrary set of the given density has such a coefficient.

[write] **6 October 2026; sources 06–10.** Corollary I.24.3 has exactly source 01’s length (I.6.1) and source 05’s increment $\delta\eta/(4\delta - \eta)$ (Theorem I.17.4), so source 01’s Theorem I.6.1 gives an increment at least 3/2 times as large at the same length; along the frontier of Theorem I.24.2 the comparison holds for $\gamma \leq 1/2$ and changes at $\gamma = 3/(4 - \kappa)$ (proof in the note after Table 3). The repository’s proof-level pair of Corollary 2.5 is `cor25_large_scale` (`Proofs02DensityIncrement.lean`); not formalized.




Figure I.24.1: The proved density–length tradeoff, evaluated at fixed $N = 10^8$, $\delta = 0.1$, and normalized correlation $\eta = 0.01$. The vertical axis divides the guaranteed increment by η . These points are evaluations of formulas, conditional on the correlation hypothesis, rather than measurements of extremal sets.

I.24.2 Using the exact variance in the U^2 inverse step

Proposition I.24.4. *Let A be a set of density $\delta \in (0, 1)$ in a finite abelian group, and $f = \mathbf{1}_A - \delta$. Then*

$$\max_{\chi \neq 1} |\widehat{f}(\chi)| \geq \frac{\|f\|_{U^2}^2}{\sqrt{\delta(1-\delta)}}.$$

More generally, a mean-zero function with $\|f\|_2^2 = V > 0$ has a nontrivial Fourier coefficient of modulus at least $\|f\|_{U^2}^2 / \sqrt{V}$.

Proof. Fourier orthogonality gives $\|f\|_{U^2}^4 = \sum_{\chi} |\widehat{f}(\chi)|^4$ and Parseval gives $\sum_{\chi} |\widehat{f}(\chi)|^2 = V$. The trivial coefficient vanishes. Bound one of the two quadratic factors in each fourth power by the largest coefficient squared and divide by V . $\square$

In a cyclic group, this may be substituted directly into Corollary I.24.3. It is a variance-sensitive version of the standard U^2 inverse argument, not a replacement for the higher-order inverse theorem. The proof also explains why an apparent improvement obtained solely by changing normalized to unnormalized Fourier coefficients would be spurious.

I.25 Conventions and normalization

Source 13, delivered Section 2.

All expectations are normalized. For a nonempty finite set X , $\mathbb{E}_X f = |X|^{-1} \sum_{x \in X} f(x)$. Write $\mathbb{T} = \mathbb{R}/\mathbb{Z}$ and $e(t) = \exp(2\pi it)$. A proper integer arithmetic progression in $I_r = \{0, \dots, r-1\}$ has the form

$$P = \{u, u+d, \dots, u+(\ell-1)d\}, \quad d \geq 1,$$

with every displayed integer in I_r . It has exactly ℓ elements. Singletons are allowed. For modular progressions, injectivity of the index parametrization is a separate condition. The repository's predicate `ModAP.IsProper` is exactly equality of carrier cardinality and declared length. That predicate alone should not be confused with the existence of a nonwrapping integer lift.

Definition I.25.1 (Arc diameter). For a finite subset $S \subset \mathbb{T}$, let $\text{diam}_{\text{arc}}(S)$ be the length of the shortest closed circular arc containing it. For $S \subset \mathbb{Z}/N\mathbb{Z}$, its modular diameter is the least integer $s \geq 0$ such that S lies in a set $\{u, u+1, \dots, u+s\}$ modulo N . On identifying $\mathbb{Z}/N\mathbb{Z}$ with $(1/N)\mathbb{Z}/\mathbb{Z}$, these two diameters differ by the factor N .

Arc diameter is stronger than merely bounding all pairwise circular distances. This distinction matters in the obstruction proof. If phases lie in an arc of length $\varepsilon \leq 1/2$, their complex values lie in an arc of angle $2\pi\varepsilon$ and have chord diameter at most $2\sin(\pi\varepsilon)$.

For a finite abelian group G , the normalized Fourier transform is

$$\hat{f}(\chi) = \mathbb{E}_{x \in G} f(x) \overline{\chi(x)}.$$

Thus Parseval reads $\mathbb{E} f \bar{g} = \sum_{\chi} \hat{f}(\chi) \overline{\hat{g}(\chi)}$. Define

$$\|f\|_{U^d}^{2^d} = \mathbb{E}_{x, h_1, \dots, h_d \in G} \prod_{\omega \in \{0,1\}^d} \mathcal{C}^{|\omega|} f(x + \omega \cdot h), \quad (\text{I.25.1})$$

where $\mathcal{C}$ denotes complex conjugation. In particular,

$$\|f\|_{U^2}^4 = \sum_{\chi} |\hat{f}(\chi)|^4.$$

For $d \geq 2$ this is a norm; its triangle inequality and the cube Gowers–Cauchy–Schwarz inequality are standard [1, 7]. We recall the latter in the precise form used below:

$$\left| \mathbb{E}_{x, h} \prod_{\omega} \mathcal{C}^{|\omega|} f_{\omega}(x + \omega \cdot h) \right| \leq \prod_{\omega} \|f_{\omega}\|_{U^d}. \quad (\text{I.25.2})$$

Repeated Cauchy–Schwarz, one coordinate at a time, proves this inequality: duplicate a coordinate, eliminate factors independent of it, and continue until each surviving factor appears in its own full cube average. All averages are finite, so no convergence interchange is involved.

The paper calls f α -uniform of degree s when

$$\|f\|_{U^{s+1}}^{2^{s+1}} \leq \alpha.$$

Its degree parameter is therefore one less than the modern norm index. The Lean definition `gowersNorm d f` is unnormalized and equals

$$|G|^{(d+1)/2^d} \|f\|_{U^d}$$

in the cyclic setting. These factors must be restored when inserting any result into that development.

I.26 Simultaneous linear phases: finite partitions and optimal powers

Source 13, delivered Section 3.

Let $\phi_i(n) = \alpha_i n + \beta_i \in \mathbb{T}$, $1 \leq i \leq q$. Constants β_i have no effect on image diameter. We allow different tolerances $0 < \varepsilon_i \leq 1$, and write $E_0 = \prod_i \varepsilon_i$. The one-phase approximation and density-increment method is classical; for a modern presentation see Zhao's notes [52]. We keep the whole frequency vector in one application of the pigeonhole principle.

Lemma I.26.1 (Rectangular simultaneous approximation). *For positive integers $Q_1, \dots, Q_q$, put $D = \prod_i Q_i$. There exists an integer $1 \leq d \leq D$ such that*

$$\|d\alpha_i\|_{\mathbb{T}} < Q_i^{-1} \quad (1 \leq i \leq q).$$

Proof. Divide $[0, 1)^q$ into the D half-open boxes of side lengths $1/Q_i$. Two of the $D + 1$ points $n(\alpha_1, \dots, \alpha_q)$, $0 \leq n \leq D$, belong to the same box. Subtract their indices. Each coordinate difference has a real representative of absolute value less than $1/Q_i$. $\square$

Theorem I.26.2 (Complete partition with flexible lengths). *Let $L \geq 2$ be an integer and set*

$$Q_i = \left\lceil \frac{2L - 2}{\varepsilon_i} \right\rceil, \quad D = \prod_i Q_i.$$

If $r \geq LD$, then I_r has a partition into proper arithmetic progressions, all with one common positive difference $d \leq D$, satisfying

$$L \leq |P| \leq 2L - 1, \quad \text{diam}_{\text{arc}} \phi_i(P) \leq \varepsilon_i \quad \text{for every } i.$$

For modular phases and $r \leq N$, the image-diameter bound is $\varepsilon_i N$ in the integral convention of section I.25.

Proof. Choose d from lemma I.26.1. Each residue chain modulo d inside I_r has at least $\lfloor r/d \rfloor \geq L$ elements. Write its length as $mL + s$, where $m \geq 1$ and $0 \leq s < L$. Split the chain consecutively into one piece of length $L + s$ and $m - 1$ pieces of length L . All pieces have lengths in the required interval.

Choose a real lift θ_i of $d\alpha_i$ with $|\theta_i| < 1/Q_i$. A cell of length ℓ has phase values obtained by reducing a real arithmetic progression of span $(\ell - 1)|\theta_i|$. Its span is less than $(2L - 2)/Q_i \leq \varepsilon_i$. For a modular phase, the same lift can be chosen as an integer representative divided by N ; the real span times N is then integral, so there is no extra rounding loss in the final bound. $\square$

Corollary I.26.3 (An explicit scale for every r). *There is such a partition with minimum length*

$$L_* = \max \left\{ 1, \left\lceil \left(\frac{rE_0}{2^q} \right)^{1/(q+1)} \right\rceil \right\}. \quad (\text{I.26.1})$$

If $L_ = 1$, the partition may consist of singletons. The number M of cells satisfies $r/(2L_* - 1) \leq M \leq r/L_*$.*

Proof. For $L_* \geq 2$, use $Q_i \leq (2L_* - 1)/\varepsilon_i \leq 2L_*/\varepsilon_i$. Then $L_* \prod_i Q_i \leq 2^q L_*^{q+1}/E_0 \leq r$, and theorem I.26.2 applies. The bounds on M follow by summing the cell lengths. $\square$

[write] **6 October 2026; sources 11–14.** *Credit.* Theorem I.26.2 with Corollary I.26.3 is the simultaneous phase-flat partition of source 01’s Theorem I.5.1 and source 05’s Theorem I.18.2 (the same common-difference construction; L_* is their L), proved independently here.

I.26.1 A matching obstruction

Lemma I.26.4 (A progression in a short arc has a small lifted step). *Suppose $L \geq 2$, $0 < \varepsilon < 1/2$, and $c, c + \theta, \dots, c + (L - 1)\theta$ lie in a circular arc of length ε . Then*

$$\|\theta\|_{\mathbb{T}} \leq \frac{\varepsilon}{L - 1}.$$

Proof. Lift all points into a single real interval of length ε . Every consecutive difference lies in $[-\varepsilon, \varepsilon]$ and is congruent to θ modulo 1. Two such differences differ by an integer of absolute value less than 1, so they coincide. If their common value is t , then $|(L - 1)t| \leq \varepsilon$. $\square$

Theorem I.26.5 (Worst-case obstruction). *Let $r \geq L \geq 2$, $0 < \varepsilon_i < 1/2$, and*

$$D_L = \left\lfloor \frac{r - 1}{L - 1} \right\rfloor.$$

If

$$D_L \prod_{i=1}^q \frac{2\varepsilon_i}{L - 1} < 1, \tag{I.26.2}$$

there exist q linear torus phases for which no length- L progression in I_r satisfies all the prescribed diameter bounds. It is sufficient that

$$(L - 1)^{q+1} > 2^q(r - 1)E_0.$$

Proof. Choose the slopes independently with uniform measure on $\mathbb{T}$. For a fixed integer $1 \leq d \leq D_L$, multiplication by d preserves uniform measure. The probability that $\|d\alpha_i\|_{\mathbb{T}} \leq \varepsilon_i/(L - 1)$ for every i is exactly the product in (I.26.2). A union bound over d leaves a positive-measure set of slopes for which no such simultaneous approximation occurs. Every length- L progression in I_r has a difference in this range, and lemma I.26.4 would force the excluded approximation. Longer good progressions are also excluded because they contain length- L subprogressions. $\square$

[write] **6 October 2026; sources 15–20.** The polynomial analogues are source 18’s Corollary I.33.5 and source 19’s Theorem I.34.7; the own-step analogue is source 16’s Corollary I.31.2.

There is a fully finite counterpart. If $N \geq r$ is prime and

$$D_L \prod_{i=1}^q \left(2 \left\lfloor \frac{\varepsilon_i N}{L - 1} \right\rfloor + 1 \right) < N^q, \tag{I.26.3}$$

then some vector of slopes in $(\mathbb{Z}/N\mathbb{Z})^q$ has the same obstruction. Indeed, every $d \leq D_L < N$ acts bijectively on $\mathbb{Z}/N\mathbb{Z}$, and the expression in parentheses counts the allowed centered residues exactly. Count the union of the bad slope vectors. For each strict torus inequality, (I.26.3) holds for all sufficiently large primes N .

Corollary I.26.6 (Sharp exponent frontier). *For fixed q , flexible complete partitions with minimum cell length of order r^a and uniform diameter of order r^{-b} have the optimal growing-cell frontier*

$$(q + 1)a + qb = 1. \quad (\text{I.26.4})$$

Here $a > 0$, $b \geq 0$, and when $b = 0$ the prescribed diameter is a fixed number strictly between 0 and $1/2$. In particular, equal length and reciprocal-diameter powers have optimal exponent $a = b = 1/(2q + 1)$.

Proof. Insert $E_0 = r^{-qb}$ into (I.26.1) for achievability. If $(q + 1)a + qb > 1$ and $a \leq 1$, take L of order r^a in theorem I.26.5; its sufficient inequality eventually holds for any fixed positive leading constant whenever $L \leq r$. If $L > r$, the desired progression is already impossible. In particular $a > 1$ is excluded immediately by $|P| \leq r$. At $b = 0$, insert a fixed tolerance below $1/2$ rather than the vacuous tolerance 1. $\square$

More precisely, for tolerances $\varepsilon_i < 1/2$, set $T_0 = (rE_0)^{1/(q+1)}$. The constructive leading constant is $2^{-q/(q+1)}$. Suitable slopes have no good progression longer than $2^{q/(q+1)}T_0 + O(1)$. The ratio of these leading constants is $2^{2q/(q+1)} < 4$. Thus complete coverage costs no additional power compared with finding just one good progression.

[write] **6 October 2026; sources 11–14.** *Credit and advance.* The exponent frontier $1/(q + 1)$ is already implied by source 05's Theorem I.18.5 (equal widths, with an explicit badly approximable vector and uniformly in N), and source 01's Theorem I.7.2 is the case of one phase with a sharper constant. Theorem I.26.5 adds independent tolerances $\varepsilon_i < 1/2$, which source 05's remark after its theorem leaves open, a leading-constant ratio below 4, and the finite prime-modulus certificate (I.26.3); this advances item I.8. Only the flexible-length frontier is claimed optimal; the prescribed-length exponent $1/(3q + 2)$ below is sufficient only (the source's own non-claim).

I.26.2 Prescribing two consecutive lengths

The exact output format of Lemma 5.9 requires a different integer decomposition. This distinction is quantitative and must be kept visible.

Lemma I.26.7 (Two consecutive allowable lengths). *If $v \geq 3$ and $b \geq (v - 1)(v - 2)$, then b is a sum of positive integers each equal to $v - 1$ or v .*

Proof. Put $m = v - 1$ and write $b = sm + t$, $0 \leq t < m$. Since $s \geq m - 1 \geq t$, we have $b = (s - t)m + t(m + 1)$. $\square$

Theorem I.26.8 (Prescribed-length simultaneous partition). *For $v \geq 3$, define*

$$B_v = (v - 1)(v - 2), \quad Q_i = \left\lceil \frac{v - 1}{\varepsilon_i} \right\rceil.$$

If $r \geq B_v \prod_i Q_i$, there is a complete common-difference partition of I_r into proper progressions of lengths $v - 1$ or v , with all phase-image diameters at most ε_i . For $v = 1$ or $v = 2$, singletons suffice.

Proof. Choose the common difference by lemma I.26.1. Each residue chain has at least B_v elements; decompose its length using lemma I.26.7, and use consecutive blocks. A block has at most $v - 1$ steps, so its i th lifted phase span is at most $(v - 1)/Q_i \leq \varepsilon_i$. $\square$

Since $Q_i \leq v/\varepsilon_i$ and $B_v \leq v^2$, a convenient sufficient condition is

$$1 \leq v \leq (rE_0)^{1/(q+2)}. \quad (\text{I.26.5})$$

Corollary I.26.9 (A direct strengthened linear form of Lemma 5.9). *Let $q \geq 1$, $a_q = 1/(3q+2)$, and let $\phi_1, \dots, \phi_q : \mathbb{Z}/N\mathbb{Z} \rightarrow \mathbb{Z}/N\mathbb{Z}$ be affine maps. For every $1 \leq r \leq N$ and integer $1 \leq v \leq r^{a_q}$, there is a complete partition of I_r into nonempty proper integer progressions of lengths $v-1$ or v such that*

$$\text{diam } \phi_i(P) \leq r^{-2a_q} N \quad \text{for every } i, P.$$

Proof. Take $\varepsilon_i = r^{-2a_q}$. Then $(rE_0)^{1/(q+2)} = r^{(1-2qa_q)/(q+2)} = r^{a_q}$, and apply (I.26.5). If $v = 1$, the only allowed nonempty size is 1, so there is no empty-cell issue. $\square$

[write] **6 October 2026; sources 15–20.** Sources 18 and 19 extend this to polynomial phases, conditionally on Maynard’s theorem (source 18’s Theorem I.32.2, which in degree one has this exponent $1/(3q+2)$; source 19’s Theorem I.34.3); this answers source 13’s Question 3 (item I.1). Source 16’s Section I.30 treats the different own-step problem for affine frequencies.

[write] **6 October 2026; sources 21–27.** *The exponent is optimal.* Source 23’s Proposition I.38.1 shows that no partition into lengths $v-1, v$ reaches a region beyond $(q+2)a + qb \leq 1$, even with cell-dependent differences; at Gowers’s $b = 2a$ this is $a \leq 1/(3q+2)$. Its Theorem I.36.1 gives the threshold with leading constant one (for lengths $L, L+1$, Corollary I.37.1). This answers source 13’s Question 1 (Research question V.54.1; item I.12).

At degree $k = 1$, the repository’s polynomial constant is $K = (k!)^2 2^{(k+1)^2} = 16$. Its target exponent is $a_q^G = 1/(2 \cdot 16^q)$ and its diameter exponent is $2a_q^G$. Since $a_q > a_q^G$, corollary I.26.9 simultaneously permits every old target length and gives a smaller image diameter. It also applies at every finite size. The comparison is strictly with the linear specialization; no claim is made for degree $k \geq 2$. [6 October 2026; sources 11–14] *A further formal baseline:* the development also has the one-phase linear partition `section5_strong_linear_partition` (`Proofs05LinearPartition.lean`, line 86 at HEAD `e1db00e99`), with lengths $v-1$ or v for $1 \leq v \leq r^{1/16}$, diameter $r^{-1/8}N$ and $r \geq 65536$. For $q = 1$ the exponent $1/5$ here still beats its $1/16$, with no size threshold; for $q \geq 2$ the formal development has no simultaneous linear statement.

q	Lemma 5.9, $k = 1$	Corollary I.26.9	Flexible lengths, $b = 2a$
1	1/32	1/5	1/4
2	1/512	1/8	1/7
3	1/8192	1/11	1/10
4	1/131072	1/14	1/13

The flexible frontier is proved optimal; the prescribed-length exponent $1/(3q+2)$ is a sufficient exponent. Its extra denominator term comes from the conductor $(v-1)(v-2)$, rather than from the simultaneous approximation step. Removing that loss is a concrete further problem.

I.26.3 Exact lengths with an exceptional set

For $v \geq 2$, choose $Q_i = \lceil (v-1)/\varepsilon_i \rceil$ and a difference $d \leq D = \prod_i Q_i$. Divide each residue chain into blocks of exactly v and discard its final remainder. Fewer than v points are discarded per chain, so the exceptional set has size at most $D(v-1)$. Hence

$$v^{q+1} \leq \vartheta r E_0$$

is sufficient to retain at least $(1 - \vartheta)r$ points in equal-length good progressions. This recovers the flexible exponent at the cost of a measured exceptional set. Such a result can be used downstream only if the exceptional mass is charged to an available budget.

I.26.4 Fewer generators, with coefficient size controlled

Suppose $\alpha_i = \sum_{j=1}^h m_{ij}\gamma_j$ in $\mathbb{T}$, where $m_{ij} \in \mathbb{Z}$. If positive integers Q_j satisfy

$$(2L - 2) \sum_j \frac{|m_{ij}|}{Q_j} \leq \varepsilon_i \quad \text{for all } i, \quad L \prod_j Q_j \leq r,$$

then the flexible partition theorem applies after approximating only the h generators. Indeed, the lift of $d\alpha_i$ is the corresponding integer combination of the small generator lifts. If $H = \max_i \sum_j |m_{ij}| \geq 1$ and all tolerances equal ε , we obtain

$$L = \max \left\{ 1, \left\lceil \left(\frac{r(\varepsilon/H)^h}{2^h} \right)^{1/(h+1)} \right\rceil \right\}. \quad (\text{I.26.6})$$

The power depends on h , and the coefficient cost is explicit.

An abstract rank without a height bound would be misleading: every modular slope is an integer multiple of $1/N$, but those integers can be of order N . That description does not give a free one-dimensional improvement. For a specified matrix, one can choose positive budgets t_j satisfying $\sum_j |m_{ij}|t_j \leq \varepsilon_i$ and maximize $\sum_j \log t_j$. This finite concave maximization gives the best rectangular budget for that presentation. It does not claim that the presentation or a rectangle is globally optimal.

I.27 Sharp selection from partition imbalance

Source 13, delivered Section 4.

Let (X, μ) be a finite probability space, with positive atom weights, and let $P_1, \dots, P_M$ form a partition of X into nonempty cells. Write

$$w_j = \mu(P_j), \quad x_j = \mathbb{E}(f \mid P_j), \quad D = \sum_j w_j |x_j|.$$

Assume throughout this section that $-a \leq f \leq b$, where $a, b > 0$, and $\mathbb{E}f = 0$. For a balanced indicator, $a = \delta$ and $b = 1 - \delta$. In a uniform space of N points, $w_j = |P_j|/N$.

Let $P_+ = \sum_j w_j (x_j)_+$. The zero-mean condition gives $P_+ = D/2$. The mass of negative-average cells is at least P_+/a , and the mass of positive-average cells is at least P_+/b . Consequently

$$0 \leq P_+ \leq \frac{ab}{a+b}. \quad (\text{I.27.1})$$

This negative-support constraint is precisely the information absent from the usual two-cutoff proof of Lemma 5.15.

Theorem I.27.1 (The exact high-average mass bound). *Suppose $D \geq 2p > 0$. For $0 \leq t < b$, define*

$$R_p(t) = \frac{p(1 + t/a) - t}{b - t}. \quad (\text{I.27.2})$$

The total mass of cells with $x_j > t$ is at least $\max\{0, R_p(t)\}$. If an exceptional union of cells has mass at most e , the nonexceptional high-average mass is at least $\max\{0, R_p(t) - e\}$. These mass bounds are sharp in the finite weighted category.

Proof. Let h be the high-average mass. Positive cells outside that union have mass at most $1 - h - P_+/a$, so

$$P_+ \leq bh + t(1 - h - P_+/a).$$

Rearranging and using $P_+ \geq p$ proves $h \geq R_p(t)$. Removing mass at most e gives the exceptional version.

For sharpness when $R_p(t) > 0$, use three cell types: average b on mass $h = R_p(t)$, average $-a$ on mass p/a , and average t on the remaining mass. The defining equation for h makes the total positive discrepancy p and the mean zero. All weights are nonnegative by (I.27.1). The high-average mass is exactly h ; at $t = 0$ the remainder has zero average. If $R_p(t) \leq 0$, use averages $-a$ and $ap/(a - p) \leq t$, with negative-cell mass p/a . There is then no high-average cell. For an exceptional budget, split the b -cell if necessary and mark the desired portion exceptional. If that budget exceeds h , all high cells can be marked exceptional. Rational approximations give the corresponding limiting sharpness for uniform finite spaces. $\square$

[write] **6 October 2026; sources 11–14.** *Credit.* The mass bound $R_p(t)$ is the sharp one-sided bound of Part I, now proved in seven sources: it is source 02's $Q_{a,b}(p, t)$ (Theorem I.9.1) and source 03's Theorem I.14.1, with the extremizers $\{-a, t, b\}$; see the note after Theorem I.4.1. The corollary below is the $M - 1$ corollary of sources 02 and 03 (Corollary I.9.3); source 06's exact fixed- M optimum (Theorem I.21.1) is at least as large, strictly for $\alpha < 1$ (note after Corollary I.9.3), and it answers this source's own question on the largest cell for a fixed number of cells in the weighted model (item I.6). The drift form (I.27.5) is source 06's signed-mean version (Section I.21).

Corollary I.27.2 (Size and density together). *If $R_p(t) > 0$, there is a cell with $x_j > t$ and*

$$\mu(P_j) \geq \frac{R_p(t)}{M - 1}.$$

With M_{good} nonexceptional cells, a positive bound $R_p(t) - e$ instead gives a good high-average cell of mass at least $(R_p(t) - e)/M_{\text{good}}$.

Proof. Positive discrepancy and zero mean force at least one negative-average cell, so at most $M - 1$ cells are high. Average the mass bound over those cells. The exceptional statement uses the safe total count M_{good} . $\square$

The high-average *mass* frontier is sharp. Dividing by a number of cells is a useful consequence; we do not assert that this division solves every extremal problem with a prescribed value of M .

Corollary I.27.3 (Discarded mass and a maximal average). *If a union of cells of mass at most $s < p/b$ is discarded, some retained cell has average at least*

$$\frac{p - bs}{1 - s - p/a}. \quad (\text{I.27.3})$$

In particular,

$$\max_j x_j \geq \frac{ap}{a - p}. \quad (\text{I.27.4})$$

Both bounds are sharp for finite weighted partitions.

Proof. The denominator is positive because $1 - p/a \geq p/b > s$. For every t below the value in (I.27.3), direct rearrangement gives $R_p(t) > s$. Thus some retained cell has average greater than t . Take the maximum of the finitely many retained averages and let t increase to the displayed value. For equality, use a discarded cell of average b and mass s , a retained cell of average $-a$ and mass p/a , and a retained cell of average (I.27.3) on the remaining mass. $\square$

For the original symmetric range $[-1, 1]$, set $a = b = 1$ and $p = \alpha/2$. At $t = \alpha/4$, corollary I.27.2 gives

$$|P_j| \geq \frac{\alpha(2 + \alpha)}{2(4 - \alpha)(M - 1)}N, \quad \mathbb{E}_{P_j} f > \frac{\alpha}{4}.$$

Lemma 5.15 gives $\alpha N/(4M)$ at the same density threshold. The larger advantage of the general theorem is that a and b can retain the actual ambient density, and t remains a free parameter.

I.27.1 Restriction with a nonzero mean

Restricting an ambient set can create a small mean drift. It need not be hidden in an absolute error. If $m = \mathbb{E}f$, put

$$P_+ = (D + m)/2, \quad P_- = (D - m)/2.$$

The same proof gives

$$\mu\{x_j > t\} \geq \max \left\{ 0, \frac{P_+ + tP_-/a - t}{b - t} \right\}. \quad (\text{I.27.5})$$

After discarding mass $s < P_+/b$, a retained average is at least $(P_+ - bs)/(1 - s - P_-/a)$. The three-type construction from corollary I.27.3, with negative mass P_-/a , proves sharpness. This formulation preserves the sign of the drift rather than paying for its absolute value twice.

I.28 Real phase removal and a joint extremal theorem

Source 13, delivered Section 5.

The real-valuedness assumption in Lemma 5.14 permits a stronger estimate than a direct triangle-inequality subtraction. Let $z : X \rightarrow \mathbb{C}$ satisfy $|z| = 1$. For one cell, use unnormalized integrals and put

$$c = \left| \int_P f z d\mu \right|, \quad d = \left| \int_P f d\mu \right|, \quad v = \int_P |f| d\mu.$$

Theorem I.28.1 (An exact cell envelope). *If f is real and the phase values on P have chord diameter at most $2s$, where $0 \leq s \leq 1$, then*

$$c^2 \leq (1 - s^2)d^2 + s^2v^2. \quad (\text{I.28.1})$$

For every permissible pair of positive and negative masses, equality can be attained.

Proof. Let u_+ and u_- be the positive and negative integrals of f on P . Then $d = |u_+ - u_-|$ and $v = u_+ + u_-$. For any pair of phase values, the chord condition gives $\Re(z_x \bar{z}_y) \geq 1 - 2s^2$. Expanding the squared correlation, bounding the two same-sign terms by u_+^2 and u_-^2 , and using this lower bound on the subtracted cross term gives

$$c^2 \leq u_+^2 + u_-^2 - 2u_+u_-(1 - 2s^2) = d^2 + 4s^2u_+u_-.$$

This is (I.28.1). Equality holds if all positive mass has one unit phase and all negative mass another unit phase at chord distance $2s$. $\square$

Corollary I.28.2 (The optimal affine aggregate envelope). *If every cell has chord diameter at most $2s$, define*

$$C = \sum_j \left| \int_{P_j} f z d\mu \right|, \quad V = \int_X |f| d\mu.$$

Then

$$C \leq sV + (1-s)D. \quad (\text{I.28.2})$$

For $s < 1$, this gives $D \geq \max\{0, (C - sV)/(1-s)\}$. The coefficients in (I.28.2) are optimal using only these data.

Proof. For $0 \leq d \leq v$, the square of $sv + (1-s)d$ exceeds $(1-s^2)d^2 + s^2v^2$ by $2s(1-s)d(v-d) \geq 0$. Apply theorem I.28.1 and sum. Sign-pure cells attain $d = v, c = v$; sign-balanced cells at the two phase endpoints attain $d = 0, c = sv$. Mixtures attain every point on the asserted affine envelope. The signs of sign-pure cells can be paired if zero mean is also required. $\square$

[write] **6 October 2026; sources 11–14.** *Credit.* Theorem I.28.1 is the one-cell envelope of source 01's Lemma I.1.2 and source 02's Theorem I.8.1, and this corollary is source 01's linear law (Lemma I.1.1) and source 02's Corollary I.8.2; the complex two-atom example below plays the role of source 02's complex counterexample. Proved independently here; printed in this source's form.

For a balanced set, $V = 2\delta(1-\delta)$ exactly. For general mean-zero $f \in [-a, b]$, $V \leq 2ab/(a+b)$. Thus a fixed target $D \geq \vartheta\alpha$ from $C \geq \alpha$, $0 < \vartheta < 1$, is ensured by the single global choice

$$s \leq \frac{\alpha(1-\vartheta)}{V - \vartheta\alpha}. \quad (\text{I.28.3})$$

Here $V \geq C \geq \alpha$, so the denominator is positive. For $\vartheta = 1/2$, the condition is $s \leq \alpha/(2V - \alpha)$.

This separation is useful in Lemma 5.14. First construct a refinement with a globally chosen phase tolerance; then apply (I.28.2). A cell-count constant obtained by applying a polynomial partition theorem separately to local correlations α_j is not automatically uniform as $\alpha_j \rightarrow 0$. A global tolerance, or a global cutoff with its discarded correlation charged explicitly, prevents that hidden dependency. The repository's full-partition proof is a relevant insertion point [68].

The real hypothesis is essential. With two equal-weight atoms, take phases $e^{i\theta}, e^{-i\theta}$ and function values $e^{-i\theta}, e^{i\theta}$, $0 < \theta < \pi/2$. Then $C = V = 1$, $D = \cos \theta$, and $s = \sin \theta$, so the right side of (I.28.2) is strictly less than 1. The result applies to real phase-removal and balanced-set steps, rather than to an arbitrary complex-valued input.

I.28.1 Solving phase geometry and density selection together

Composing corollary I.28.2 with theorem I.27.1 is valid, but can lose information: an extremizer for the first bound need not be an extremizer for the second. The following theorem optimizes both constraints at once.

For $-a \leq x \leq b$, define

$$v(x) = \frac{2ab + (b-a)x}{a+b}, \quad (\text{I.28.4})$$

$$h_s(x) = \sqrt{(1-s^2)x^2 + s^2v(x)^2}. \quad (\text{I.28.5})$$

For $0 \leq t < b$ and $0 \leq s < 1$, put

$$H_s(t) = \frac{a(t + h_s(t))}{a + t}, \quad K_s(t) = \frac{(a + b)(v(t) - h_s(t))}{a + t}. \quad (\text{I.28.6})$$

The number $K_s(t)$ is strictly positive: $v(t) > t \geq 0$ for $t < b$, and $s < 1$.

Theorem I.28.3 (Sharp joint phase, density, and exceptional-mass frontier). *Let $-a \leq f \leq b$, $\mathbb{E}f = 0$, and let C be the aggregate cell correlation above. Suppose phase chord diameter is at most $2s$ on every nonexceptional cell, where $0 \leq s < 1$. Let e be exceptional mass, and q_t the mass of nonexceptional cells with average greater than t . Then*

$$C \leq H_s(t) + K_s(t)(e + q_t). \quad (\text{I.28.7})$$

In particular, if $C \geq \alpha$ and $e \leq e_0$,

$$q_t \geq \max \left\{ 0, \frac{\alpha - H_s(t)}{K_s(t)} - e_0 \right\}. \quad (\text{I.28.8})$$

This is the exact lower frontier for high-average mass in the finite weighted problem with the stated information.

[write] **6 October 2026; sources 11–14.** *Credit.* This joint frontier is the content of source 01's Theorems I.2.1, I.2.3 and I.3.1 (the short-arc density tail with exceptional mass and the optimal forced maximum); the intake found (H_s, K_s) equal to source 01's $(B_{\delta,s}, \Gamma_{\delta,s})$ numerically. Source 01 charges exceptions by occupied mass, which is stronger for indicators; “exact” here refers to the mass budget stated.

Proof. The chord joining $(-a, a)$ and (b, b) majorizes absolute value, so $|f| \leq v(f)$ pointwise. Therefore a cell of average x has conditional absolute integral at most $v(x)$. By theorem I.28.1, a nonexceptional cell has conditional correlation at most $h_s(x)$.

The function h_s is convex, being the Euclidean norm of the affine vector $(\sqrt{1 - s^2}x, sv(x))$. Let ℓ_t be the line through $(-a, a)$ and $(t, h_s(t))$. Convexity gives $h_s(x) \leq \ell_t(x)$ on $[-a, t]$. The affine function $v(x) - \ell_t(x)$ vanishes at $-a$, has nonnegative slope, and at b equals $K_s(t)$. Thus every exceptional or high-average cell has conditional correlation at most

$$v(x) \leq \ell_t(x) + K_s(t),$$

whereas every other cell has conditional correlation at most $\ell_t(x)$. Average these inequalities. Since $\mathbb{E}f = 0$, the affine terms sum to $\ell_t(0) = H_s(t)$, proving (I.28.7) and (I.28.8).

For sharpness, put $A_0 = a + b$ and choose $0 \leq u \leq a/A_0$. Use three cell types with masses and averages

Type	Mass	Average
High	u	b
Threshold	$(a - A_0 u)/(a + t)$	t
Negative	$(t + (b - t)u)/(a + t)$	$-a$

The masses sum to 1 and the mean is zero. On the threshold cell, take values b and $-a$ in proportions $(a+t)/A_0$ and $(b-t)/A_0$, and assign them unit phases at chord distance $2s$. Its conditional correlation is exactly $h_s(t)$. Give the other two cells constant phases. Direct substitution yields

$$C = H_s(t) + K_s(t)u.$$

Any mass $e \leq u$ of the high cell may be split off as exceptional, leaving $q_t = u - e$. The identity

$$H_s(t) + \frac{a}{a+b} K_s(t) = \frac{2ab}{a+b} \quad (\text{I.28.9})$$

shows that these constructions cover every α between $H_s(t)$ and the maximal possible correlation $2ab/(a+b)$. If $\alpha \leq H_s(t)$, take $u = 0$; if the exceptional budget covers the required u , declare the whole high cell exceptional. In either case the zero lower bound is attainable. This proves the full sharpness statement. $\square$

The extremizers use only the endpoint values $-a, b$. For $a = \delta$, $b = 1 - \delta$, those are exactly the values of a balanced indicator. Finite weighted realizations are exact, and uniform finite realizations approximate arbitrary real weights. However, an additional requirement that the phase be one fixed low-degree polynomial can restrict the extremizers; the theorem does not assert optimality under that extra structure.

I.28.2 An exact maximum-average guarantee

With no exceptional cells, H_s is continuous and strictly increasing on $[0, b]$, with

$$H_s(0) = s \frac{2ab}{a+b}, \quad H_s(b) = \frac{2ab}{a+b}.$$

For $0 < s < 1$, strict convexity of h_s makes the secant slopes from $(-a, a)$ strictly increase. For $s = 0$, the formula is $H_0(t) = 2at/(a+t)$, which is strictly increasing directly. Thus, whenever

$$s \frac{2ab}{a+b} < \alpha \leq \frac{2ab}{a+b},$$

the exact least possible maximum cell average is the unique $t_* \in (0, b]$ satisfying $H_s(t_*) = \alpha$. A two-cell construction with averages $-a, t_*$ attains it. When a partition already guarantees a minimum cell length, this gives a density increment with no division by the cell count.

For example, $\delta = 0.2$, $\alpha = 0.04$, and $s = 0.1$ give $V = 0.32$. Separate affine phase removal and selection guarantee a maximum average of $0.00454545\dots$; the joint theorem gives $0.00814479\dots$. At $t = 0.005$, it forces high-average mass at least $0.00228653\dots$, whereas the separate composition gives no positive bound. These are evaluations of explicit formulas, rather than empirical estimates of an unknown extremum.

I.29 A density-sensitive Fourier increment

Source 13, delivered Section 6.

The partition and selection results combine into an explicit finite corollary. It improves the actual strengthened conclusion already visible in the corrected proof of Corollary 2.5, rather than merely restating that correction.

Theorem I.29.1 (Fourier correlation yields a longer dense progression). *Let $A \subset \mathbb{Z}/N\mathbb{Z}$, $\delta = |A|/N$, and suppose some nonzero frequency r satisfies*

$$\left| \sum_{x \in A} e(-rx/N) \right| \geq \alpha N > 0.$$




Figure I.28.1: Guaranteed high-average mass at $a = 0.2$, $b = 0.8$, $\alpha = 0.04$, $s = 0.1$, with no exceptional cells. The joint frontier retains information lost by applying two sharp inequalities separately. The dotted vertical line marks the example $t = 0.005$.

There is a proper integer progression $P \subset \{0, \dots, N-1\}$ such that

$$|P| \geq \sqrt{\frac{\alpha N}{32\pi\delta(1-\delta)}} \quad (\text{I.29.1})$$

and

$$\frac{|A \cap P|}{|P|} \geq \delta + \frac{\alpha}{4 - \alpha(2-\delta)/(\delta(1-\delta))} \geq \delta + \frac{\alpha}{4}. \quad (\text{I.29.2})$$

The intersection uses the standard representative of each residue. No primality or additional size hypothesis is needed.

Proof. Set $a = \delta$, $b = 1 - \delta$. The nontrivial character sums to zero on the group, so its sum on A is the negative of its sum on the complement. Hence

$$0 < \alpha \leq \min(a, b) \leq 2ab, \quad 0 < a, b < 1.$$

Put $\rho_0 = \alpha/(4ab) \leq 1/2$ and $\varepsilon = \arcsin(\rho_0)/\pi \leq 1/6$. Apply corollary I.26.3 for one phase to rx/N on I_N . Every cell has length at least

$$L = \max\{1, \lfloor \sqrt{\varepsilon N/2} \rfloor\}$$

and phase chord diameter at most $2\rho_0$. For $f = \mathbf{1}_A - a$, one has $\mathbb{E}f = 0$ and $V = 2ab$. The aggregate cell correlation is at least α , by the triangle inequality. Therefore corollary I.28.2 gives

$$D \geq \frac{\alpha - 2ab\rho_0}{1 - \rho_0} = \frac{\alpha}{2(1 - \rho_0)}.$$

Let $p_0 = \alpha/[4(1 - \rho_0)]$. Then $p_0 \leq ab < a$, and corollary I.27.3 gives a cell with average at least

$$\frac{ap_0}{a - p_0} = \frac{\alpha}{4 - \alpha(1 + b)/(ab)}.$$

The denominator is positive, since $\alpha \leq 2ab$ gives a lower bound $2a > 0$. This increment is at most b and at least $\alpha/4$, so it is compatible with all finite and singleton cases.

Finally, $\max\{1, \lfloor x \rfloor\} \geq x/2$ for every $x \geq 0$, and $\arcsin \rho_0 \geq \rho_0$. Thus

$$L \geq \frac{1}{2} \sqrt{\frac{\varepsilon N}{2}} \geq \sqrt{\frac{\alpha N}{32\pi ab}}.$$

Every cell in the partition has this length, proving both conclusions for the selected cell. $\square$

The repository's Corollary 2.5 statement gives length $(\alpha^3 N/(128\pi))^{1/2}$ and increment $\alpha/8$. Its corrected proof already obtains length $(\alpha N/(128\pi))^{1/2}$ and increment $\alpha/4$ [67]. The additional length gain in theorem I.29.1, relative to that stronger proof output, is

$$\frac{2}{\sqrt{\delta(1-\delta)}}.$$

It comes from retaining the exact absolute mass of the balanced function, the real phase envelope, and the finite simultaneous-partition construction. A sharper increment is available by solving $H_{\rho_0}(t) = \alpha$ from theorem I.28.3; (I.29.2) has the advantage of a simple closed form.

[write] **6 October 2026; sources 11–14.** *Comparison with source 01 (proved at the third write).* Theorem I.29.1 is dominated by source 01's Theorem I.6.1 at the same length. Its partition has every cell of length $L = \max\{1, \lfloor \sqrt{N \arcsin(\rho_0)/(2\pi)} \rfloor\}$ with $\rho_0 = \alpha/(4\delta(1-\delta))$, which is the length (I.6.1) of source 01 (whose η is this α). Put $b = 1 - \delta$ and $\kappa = \alpha/(2\delta b) \in (0, 1]$. Since $2 - \delta = 1 + b$, the increment (I.29.2) is $\alpha/(4 - 2\kappa(1 + b)) = \delta b \kappa / (2 - (2 - \delta)\kappa)$, and source 01's is $I = 3\delta b \kappa / (4 - (4 - 3\delta)\kappa)$. Both denominators are positive (they are at least $\delta\kappa$ and $3\delta\kappa$), and

$$\frac{I}{\delta b \kappa / (2 - (2 - \delta)\kappa)} = \frac{3(2 - (2 - \delta)\kappa)}{4 - (4 - 3\delta)\kappa} = 1 + \frac{2(1 - \kappa)}{4 - (4 - 3\delta)\kappa},$$

because $3(2 - (2 - \delta)\kappa) - (4 - (4 - 3\delta)\kappa) = 2 - 2\kappa$. So source 01's increment is at least this one, with equality only at $\kappa = 1$, and the ratio is at most $3/2$ (the denominator is at least $4 - 4\kappa$). (Exact rational check of the identity at the write: 200,000 random (δ, κ) , ratios between 1 and 1.499997.) The “sharper increment” by solving $H_{\rho_0}(t) = \alpha$ is the route of source 01's optimal forced maximum (Theorem I.3.1). The length lower bound (I.29.1) is a factor $\sqrt{2}$ below source 05's $\sqrt{\alpha N / (16\pi\delta(1-\delta))}$; the integer length is the same as source 01's. So the abstract's and Section F.18's presentation of this theorem as a gain is a gain over the formal Corollary 2.5 only; within this report it is a second route, printed in its own form.

I.30 Direct partitions for affine frequencies and their own steps

Source 16, delivered Section 6.

I.30.1 The finite certificate theorem

Let $[R]_0 = \{0, \dots, R-1\}$, and let $\|\cdot\|_{\mathbb{T}}$ denote distance to the nearest integer. All integer progressions in this section have positive steps and are contained in $[R]_0$.

Lemma I.30.1 (Two elementary tools). *An interval with at least $L \geq 1$ integer points can be partitioned into consecutive intervals with lengths in $[L, 2L - 1]$. Also, for $\beta_i \in \mathbb{T}$ and positive integers J_i , there is $1 \leq t \leq \prod_i J_i$ with*

$$\|t\beta_i\|_{\mathbb{T}} \leq J_i^{-1} \quad \text{for every } i.$$

Proof. For the first statement, write the length as $uL + r$, $u \geq 1$, $0 \leq r < L$, and use $u - 1$ blocks of length L and one of length $L + r$. For the second, place the $\prod J_i + 1$ points $j\beta$, $0 \leq j \leq \prod J_i$, in $\prod J_i$ half-open rectangular boxes of side lengths $1/J_i$. Two share a box; subtract their indices. $\square$

Theorem I.30.2 (Certificate for an exact step-controlled partition). *Let $q \geq 1$, $L \geq 2$, $A_i, B_i \in \mathbb{T}$, and $0 < \delta_i \leq 1/2$. Put*

$$J_i = \left\lceil \frac{2}{\delta_i} \right\rceil, \quad T = \prod_{i=1}^q J_i, \quad M_0 = LT, \quad \varepsilon_i = \frac{\delta_i}{4LT^2}. \quad (6.1)$$

Suppose an integer s satisfies

$$1 \leq s \leq \left\lfloor \frac{R}{LT} \right\rfloor, \quad \|A_i s^2\|_{\mathbb{T}} \leq \varepsilon_i \quad (1 \leq i \leq q). \quad (6.2)$$

Then $[R]_0$ has an exact partition into progressions Q , with

$$L \leq |Q| \leq 2L - 1, \quad \|(A_i n + B_i)d_Q\|_{\mathbb{T}} \leq \delta_i \quad (n \in Q, 1 \leq i \leq q). \quad (6.3)$$

No primality, irrationality, or genericity hypothesis is imposed on the coefficients.

Proof. Split $[R]_0$ into residue classes modulo s . Every class has at least $\lfloor R/s \rfloor \geq LT$ points. In each class, partition its ordered index interval into blocks

$$P = \{x + sj : 0 \leq j < m'\}, \quad LT \leq m' \leq 2LT - 1,$$

using Lemma I.30.1. On one such block put $\beta_i = (A_i x + B_i)s$. The rectangular Dirichlet argument gives $1 \leq t \leq T$ with $\|t\beta_i\|_{\mathbb{T}} \leq \delta_i/2$.

Split $[m']_0$ into residue classes modulo t . Each has at least $\lfloor m'/t \rfloor \geq L$ points, so it can again be divided into blocks of lengths in $[L, 2L - 1]$. Their images in the original interval have actual step $d_Q = st$. For every $n = x + sj$ in the intermediate block,

$$(A_i n + B_i)st = t\beta_i + jt(A_i s^2) \quad \text{in } \mathbb{T}.$$

Since $j < 2LT$ and $t \leq T$,

$$\|(A_i n + B_i)st\|_{\mathbb{T}} \leq \frac{\delta_i}{2} + (2LT)T \frac{\delta_i}{4LT^2} \leq \delta_i.$$

All splits use disjoint residue classes or consecutive blocks. They therefore preserve exact coverage and disjointness, and each final cell has the required length. $\square$

The distinction from quadratic flattening is concrete: at the second recurrence step the tolerance is δ_i , rather than a tolerance divided by the final cell length. The first recurrence controls the variation of the affine frequency along that step.

I.30.2 The external recurrence theorem and the finite threshold

We use the following precise consequence of [73, Theorem 1.1]. Its proof is external to this article.

Proposition I.30.3 (Cited simultaneous quadratic recurrence). *For every $C > 21$ and $q \geq 1$, there is a constant $X_0(C, q)$ such that the following holds. If $0 < \eta_i \leq 1/100$, $X > X_0(C, q)$, and*

$$\prod_i \eta_i^{-1} \leq X^{1/C}, \quad (6.4)$$

then, for arbitrary $A_i \in \mathbb{T}$, there is an integer $1 \leq s < X$ such that $\|A_i s^2\|_{\mathbb{T}} \leq \eta_i$ for every i .

[write] **6 October 2026; sources 15–20.** *External input.* The cited preprint is not reproved here and its threshold $X_0(C, q)$ is not numerical (the source says so). The intake fetched it: its Theorem 1.1 gives $c_2 = 10.5 + 9m/(2m)^M + \varepsilon$, so $C > 21$ at degree two is accurate; arXiv lists the author as Cheuk Fung Lau. Sources 18 and 19 use the same input in Part I (Sections I.32 and I.34).

To verify the specialization, Lau's theorem has degree d and constant $c_2 = 10.5 + 9q/(2q)^M + \varepsilon$, with adjustable M, ε subject to his stated lower bound. For $d = 2$, first make $2c_2 < C$. His requirement $\prod \eta_i^{-1} \leq X^{1/(2c_2)}$ then follows from the stronger (6.4). The choices are fixed in terms of C, q and absorbed in X_0 . Maynard's published result [72] supplies a corresponding input with a degree-dependent constant and suffices for the q^{-2} nature of the conclusion below. The numerical limiting constant 42 uses Lau's stated theorem.

Corollary I.30.4 (Finite sufficient condition). *In Theorem I.30.2, set $\Delta = \prod_i \delta_i$ and $Q_0 = \lfloor R/(LT) \rfloor$. For any $C > 21$, it is sufficient that*

$$Q_0 > X_0(C, q), \quad Q_0 \geq (4LT^2)^{Cq} \Delta^{-C}. \quad (6.5)$$

Proof. The parameters ε_i in (6.1) are at most $1/100$, and

$$\prod_i \varepsilon_i^{-1} = (4LT^2)^q \Delta^{-1}.$$

Apply Proposition I.30.3 with $X = Q_0$. Its integer $s < Q_0$ is an admissible certificate. $\square$

The explicit certificate can be searched for and verified at a particular finite size. The sufficient existence condition (6.5) contains the nonnumerical constant X_0 ; these are different kinds of information.

I.30.3 The anisotropic exponent region

Theorem I.30.5 (Power-law partition region). *Fix $q \geq 1$, $C > 21$, and positive $a, b_1, \dots, b_q$. If*

$$(Cq + 1)a + (C(2q + 1) + 1) \sum_{i=1}^q b_i < 1, \quad (6.6)$$

then for all sufficiently large R , uniformly in all A_i, B_i , there is a partition satisfying (6.3) with

$$L = \lceil R^a \rceil, \quad \delta_i = R^{-b_i}.$$

In particular, every

$$0 < \gamma < \frac{1}{(q + 1)(42q + 1)} \quad (6.7)$$

is valid with $a = b_1 = \dots = b_q = \gamma$.

Proof. For large R , $J_i \leq 3/\delta_i$, hence $T \leq 3^q R^{\sum b_i}$. The exponent of R in the size requirement obtained from (6.5) is

$$a + \sum b_i + Cq \left(a + 2 \sum b_i \right) + C \sum b_i,$$

which equals the left side of (6.6). The strict inequality absorbs the fixed constants, ceiling and floor effects, and $X_0(C, q)$.

If all exponents equal γ , the displayed expression is $(q + 1)(2Cq + 1)\gamma$. For any γ as in (6.7), choose $C > 21$ sufficiently close to 21, and apply the first assertion. $\square$

For comparison, the reciprocal limiting denominators are 86, 255, 508, 845 at $q = 1, 2, 3, 4$. Gowers' Lemma 13.5 has length exponent 2^{-12q} and accuracy exponent 2^{-11q} , together with a factor $1/2$ in length [1]. The theorem above improves both powers for any fixed q , once its uniform threshold is met; it does not assert the limiting exponent with an absolute constant-one threshold.

I.30.4 Transport to cyclic groups

Corollary I.30.6 (Composite moduli are allowed in the local transfer). *Let $P = \{h_0 + d_0 n : 0 \leq n < R\} \subseteq \mathbb{Z}/N\mathbb{Z}$ be properly parameterized, and let $a_i h + b_i$ be affine frequencies. Set*

$$A_i = \frac{a_i d_0^2}{N} \pmod{1}, \quad B_i = \frac{(a_i h_0 + b_i) d_0}{N} \pmod{1}. \quad (6.8)$$

A partition supplied by Theorem I.30.2 transports to a proper partition of P , and its ambient steps satisfy

$$|(a_i h + b_i) d_Q|_N \leq \delta_i N. \quad (6.9)$$

Proof. A cell with index step st has ambient step $d_Q = d_0 st$. For $h = h_0 + d_0 n$,

$$\frac{(a_i h + b_i) d_Q}{N} = (A_i n + B_i) st \pmod{1}.$$

This proves the bound. Every output index cell is contained in the injective parameter domain of P , so properness is preserved. The formula requires no inverse of d_0 , 2, or 4, and therefore remains valid for arbitrary N . $\square$

I.30.5 Weighted extraction for the Section 13 interface

Lemma I.30.7 (Selecting a weighted cell). *Suppose $0 \leq w(n) \leq U_0$ and $\sum_{n \in [R]_0} w(n) \geq \mu U_0 R$. Every exact partition has a cell Q with $\sum_{n \in Q} w(n) \geq \mu U_0 |Q|$. For $0 \leq \tau < \mu \leq 1$, this cell has at least*

$$\frac{\mu - \tau}{1 - \tau} |Q| \tag{6.10}$$

points at which $w(n) \geq \tau U_0$.

Proof. The first statement follows by summing the contrary strict inequalities over the cells. If the high-value fraction is v , the average on Q is at most $v + (1 - v)\tau$, in units of U_0 . Solving $\mu \leq v + (1 - v)\tau$ proves (6.10). Two-valued weights at U_0 and arbitrarily below τU_0 show sharpness of the bounded-variable factor. $\square$

[write] **6 October 2026; sources 15–20.** *The Lemma 13.5 interface* (checked at HEAD 0b4890a1c). `lemma_13_5` (Sections12_13.lean, line 410) is **open**; its conclusion `IsStage135Data` (line 397) asks for a progression of length at least $|P|^{2^{-12q}}/2$, own-step frequencies at most $|P|^{-2^{-11q}} N$, and at least $\alpha^{32}|Q|/20$ critical heights. This lemma's weighted cell with $\alpha^{32}/16$ and Corollary I.30.6 (no inverse of 2, 4 or d_0) fit that shape. The formal development has moved here since the pin, independently of this source: the weighted selection now gives the same pre-trimming density $\alpha^{32}|Q|/16$ (`stage135_critical_count_of_weight`, `Proofs13QuadraticRecurrence.lean`, line 333), endpoint deletion is formalized (`stage135_of_weighted_quadratic_cell`, `Proofs13EndpointDeletion.lean`, line 146), and Lemma 13.5 is proved at an explicit sufficient scale and at large prime N with the paper's exponents (`lemma_13_5_with_scale`, `Proofs13QuadraticPartition.lean`, line 76; `lemma_13_5_large_N`, `Proofs13LargeScaleRecurrence.lean`, line 96), while the catalogue statement stays open and the module docstring of `Proofs13QuadraticRecurrence.lean` still records the two printed gaps. The polynomial exponent of Theorem I.30.5 rests on the cited preprint, with a non-numerical threshold; repository policy forbids importing it as an axiom, so only the certificate route of Theorem I.30.2 is formalizable as it stands. Nothing here is formalized; `lemma_13_5_of_candidate_family` (`Proofs13QuadraticRecurrence.lean`, line 425) is the formal conditional step it would feed.

In the notation of Gowers' Lemmas 13.4–13.5, take

$$w(h) = C(h) \mathbf{1}_{\{h \in H, G(h) \geq (1-2\eta)C(h)\}}.$$

The source input is $\sum_{h \in P} w(h) \geq \alpha^{32} N^{31} |P|/8$; the pointwise cap is N^{31} . Applying the preceding partition and weighted lemma with $\mu = \alpha^{32}/8$, $\tau = \alpha^{32}/16$, gives a cell Q having at least

$$\frac{\alpha^{32}}{16 - \alpha^{32}} |Q| \geq \frac{\alpha^{32}}{16} |Q| \tag{6.11}$$

good heights. Every such height satisfies both $C(h) \geq \alpha^{32} N^{31}/16$ and $G(h) \geq (1 - 2\eta)C(h)$. The step estimate holds throughout the full cell. In particular no endpoint deletion is used to obtain (6.11). The source comparison is the displayed denominator 20; the power improvements still require the threshold of Theorem I.30.5.

I.30.6 Effective slope dimension

The recurrence certificate only concerns A_i , not B_i . If only r of the A_i are nonzero in $\mathbb{T}$, then Corollary I.30.4 can use

$$Q_0 \geq (4LT^2)^{Cr} \left(\prod_{A_i \neq 0} \delta_i \right)^{-C}, \quad Q_0 > X_0(C, r). \quad (6.12)$$

The rectangular step t must still handle all q intercepts.

Corollary I.30.8 (Fewer varying slopes). *With $r \geq 1$ varying slopes, the sufficient power region becomes*

$$(Cr + 1)a + (2Cr + 1) \sum_i b_i + C \sum_{A_i \neq 0} b_i < 1. \quad (6.13)$$

Every balanced exponent $\gamma < ((q + 1)(42r + 1))^{-1}$ follows. If $r = 0$, the certificate $s = 1$ works whenever $R \geq LT$.

Proof. Apply the recurrence theorem only to the nonzero slopes and repeat the exponent calculation in Theorem I.30.5. At equal powers the total coefficient is $(q + 1)(2Cr + 1)$. For $r = 0$, all quadratic errors are zero. $\square$

More generally, if the A_i are integer combinations of r torus elements with row sums of absolute coefficients bounded by a fixed H , a simultaneous certificate for those r elements at accuracy $\min_i \varepsilon_i / H$ implies all the required certificates. At equal accuracies a fixed H changes thresholds, but does not change the balanced power in Corollary I.30.8. A height bound is essential for a uniform quantitative formulation.

I.31 An obstruction and a precise remaining partition problem

Source 16, delivered Section 7.

Theorem I.31.1 (Endpoint-cell obstruction). *Let $q \geq 1$, $2 \leq L \leq R$, and $0 < \delta_i < 1/4$, and set $\Delta = \prod_i \delta_i$. If*

$$\left\lfloor \frac{R-1}{L-1} \right\rfloor \frac{4^q \Delta^2}{(L-1)^q} < 1, \quad (7.1)$$

then there are coefficients $A_i, B_i \in \mathbb{T}$ for which no progression contained in $[R]_0$, containing 0, and of length at least L , satisfies (1.3). Consequently no full partition into such progressions exists.

Proof. Choose all A_i, B_i independently with Haar distribution. A progression containing 0 in $[R]_0$ begins at zero when listed in increasing order, so its step d satisfies $1 \leq d \leq \lfloor (R-1)/(L-1) \rfloor$. For this step put

$$u_i = A_i d^2, \quad v_i = B_i d.$$

The $2q$ variables remain independent and Haar distributed, since multiplication by a positive integer preserves Haar measure. The first L points can satisfy the required inequalities only if

$$\|v_i + j u_i\|_{\mathbb{T}} \leq \delta_i \quad (0 \leq j < L). \quad (7.2)$$

For a single pair (u, v) , the exact measure of (7.2) is $4\delta^2/(L-1)$. Indeed choose the unique representatives $v_j \in [-\delta, \delta]$. All adjacent differences represent u and lie in $[-2\delta, 2\delta]$, an interval of length less than one. They therefore coincide as real numbers. The sequence is a real arithmetic progression between its endpoints v_0, v_{L-1} . Conversely every pair of endpoints in $[-\delta, \delta]$ defines such a sequence, and the map to (u, v) has absolute Jacobian $1/(L-1)$. Its image has no multiplicity modulo $\mathbb{Z}^2$. The measure is $(2\delta)^2/(L-1)$.

Independence over i gives probability $4^q \Delta^2/(L-1)^q$ for each fixed step. The union bound and (7.1) leave positive probability that no candidate step works. A full partition would have a cell containing 0, proving the last assertion. $\square$

Corollary I.31.2 (Necessary power region). *For positive $a, b_1, \dots, b_q$, a coefficient-uniform full-partition theorem with $L \asymp R^a$ and $\delta_i \asymp R^{-b_i}$ can hold only if*

$$(q+1)a + 2 \sum_i b_i \leq 1. \quad (7.3)$$

In the balanced case this requires $\gamma \leq 1/(3q+1)$.

Proof. The left side of (7.1) is, up to fixed factors, at most $R^{1-(q+1)a-2\sum b_i}$. It tends to zero if the claimed necessary inequality fails. $\square$

[write] 6 October 2026; sources 15–20. *Relation to the other partition results.* This is the own-step, affine-frequency analogue of source 13's frontier $(q+1)a + qb = 1$ for linear phases and diameters (Corollary I.26.6), by the same random-slope union bound as source 13's Theorem I.26.5 and source 05's Theorem I.18.5; here q counts affine frequencies (source 13's q), not a common difference. Sources 18 and 19 prove the polynomial-phase analogues $(qk(k+1)/2+1)a + kqb \leq 1$ and $(S+1)a + Db \leq 1$ (Corollary I.33.5, Theorem I.34.7) for partitions controlling diameters, a different problem with the same recurrence input.

The obstruction can be realized with rational coefficients. For a strict instance of (7.1), the acceptable coefficients for any fixed step form a closed set, and their finite union does not fill the torus. Its complement is nonempty and open, so it contains a rational tuple with common denominator $N > 2(R-1)$. The interval $[R]_0$ then embeds properly into $\mathbb{Z}/N\mathbb{Z}$. Moreover every modular progression contained in this interval lifts to an integer progression, up to reversal: adjacent lifted differences lie in an interval of length less than N , so their common residue forces them to coincide as integers. This shows that the obstruction is also relevant to finite cyclic models.

There remains a gap between the sufficient region (6.6) and the necessary region (7.3). In particular, the balanced bounds are of orders q^{-2} and q^{-1} . The obstruction is for a *full partition*; it does not upper-bound every possible weighted extraction of one progression. That distinction is useful when deciding which improvement would have the greatest effect in Section 13.

I.32 Simultaneous polynomial partitions by descent in degree

Source 18, delivered Section 4.

Lemma 5.9 of [1] handles several polynomial phases by successively refining for one polynomial at a time. This produces an exponential dependence on the number of phases in the reciprocal partition exponent. We give a different organization of the argument: all leading coefficients are treated simultaneously, and induction is performed on the degree. The resulting combinatorial

transfer is elementary. Its analytic input is a simultaneous monomial recurrence bound, which we state explicitly rather than incorporate silently into the proof.

There are two distinct conclusions. First, established simultaneous recurrence estimates yield partition exponents with polynomial dependence on the number of phases. Second, an independent volume argument proves a necessary restriction on every universal partition theorem of this kind. The latter argument explains a remaining gap: for quadratics our general sufficient exponents have order q^{-2} , whereas the obstruction only rules out exponents larger than a constant times q^{-1} .

I.32.1 Arc diameter and the modular interpretation

Throughout this section, $\mathbb{R}/\mathbb{Z} = \mathbb{R}/\mathbb{Z}$, and $\|x\|_{\mathbb{R}/\mathbb{Z}} = \min_{m \in \mathbb{Z}} |x - m|$. For a nonempty finite set $S \subset \mathbb{R}/\mathbb{Z}$, define

$$\text{adiam}(S) = \inf\{\ell \in [0, 1] : S \subset a + [0, \ell] \pmod{1} \text{ for some } a \in \mathbb{R}/\mathbb{Z}\}.$$

This is the length of a shortest containing arc. It is stronger than merely bounding all pairwise geodesic distances. For a polynomial $F \in \mathbb{R}[X]$ and a finite set $P \subset \mathbb{Z}$, put

$$\text{osc}_P(F) = \text{adiam}\{F(x) \pmod{1} : x \in P\}.$$

If two functions have images in arcs of lengths η_1 and η_2 , then their pointwise sum has image in an arc of length at most $\eta_1 + \eta_2$ whenever this sum is at most one. We shall repeatedly use this elementary consequence of addition of real intervals.

For a modular polynomial $\phi(X) = \sum_j a_j X^j \in (\mathbb{Z}/N\mathbb{Z})[X]$, choose integer representatives of its coefficients and put $F(X) = \sum_j (a_j/N) X^j$. Then

$$\text{osc}_P(F) = \frac{1}{N} \text{diam}(\phi(P)), \quad (*)$$

where the diameter on the right is Gowers' containing-arc diameter. Indeed a shortest arc containing a finite subset of the grid $N^{-1}\mathbb{Z}/\mathbb{Z}$ may be chosen with endpoints in that grid. Thus every statement below applies directly to the modular polynomials in the repository. No primality assumption is used. On the integer parameter interval there is no restriction relating its length to N ; when that interval itself is regarded as a subset of $\mathbb{Z}/N\mathbb{Z}$, we retain $r \leq N$ so that its points remain distinct.

I.32.2 A recurrence input with an explicit interface

Fix an integer $q \geq 1$. For each $j \geq 2$, suppose that an integer-valued function $\mathcal{R}_{j,q} : (0, 1] \rightarrow \mathbb{Z}_{\geq 1}$ has the following property:

$$\begin{aligned} &\text{for every } \alpha_1, \dots, \alpha_q \in \mathbb{R}/\mathbb{Z} \text{ there exists } 1 \leq p \leq \mathcal{R}_{j,q}(\eta) \\ &\text{such that } \max_{1 \leq i \leq q} \|\alpha_i p^j\|_{\mathbb{R}/\mathbb{Z}} \leq \eta. \end{aligned} \quad (\text{I.32.1})$$

The function $\mathcal{R}_{j,q}$ is a recurrence radius. It may come from any proved recurrence theorem. This formulation separates the analytic problem of producing a common step from the combinatorial problem of turning that step into a partition covering every point.

We shall prescribe the two possible final cell lengths. The elementary integer fact that makes this possible deserves to be stated explicitly.

Lemma I.32.1 (Two integer subdivision rules). *Let $v \geq 2$ and put*

$$g(v) = \max\{1, (v-1)(v-2)\}.$$

Every integer $n \geq g(v)$ is a sum of positive integers each equal to $v-1$ or v . Also, for every integer $U \geq 1$, every integer $n \geq U$ is a sum of integers in $[U, 2U-1]$.

Proof. For the first assertion put $d = v-1$ and write $n = ad + b$ with $0 \leq b < d$. If $d \geq 2$, the hypothesis gives $a \geq d-1 \geq b$, and

$$n = (a-b)d + b(d+1).$$

The case $d = 1$ is immediate. For the second assertion write $n = aU + b$, where $a \geq 1$ and $0 \leq b < U$; use $a-1$ summands equal to U and one summand equal to $U+b$. $\square$

Define integer thresholds recursively, for $v \geq 2$ and $0 < \varepsilon \leq 1$, by

$$\mathcal{P}_{1,q}(v, \varepsilon) = g(v) \left\lceil \frac{v-1}{\varepsilon} \right\rceil^q, \quad (\text{I.32.2})$$

$$U = \mathcal{P}_{j-1,q}(v, \varepsilon/2), \quad T = \mathcal{R}_{j,q}\left(\frac{\varepsilon}{2^{j+1}U^j}\right),$$

$$\mathcal{P}_{j,q}(v, \varepsilon) = UT \quad (j \geq 2). \quad (\text{I.32.3})$$

All quantities that count steps or lengths are integers. In particular, there is no unmentioned rounding loss in this recursion.

Theorem I.32.2 (Simultaneous recurrence-to-partition transfer). *Assume (I.32.1) for $j = 2, \dots, k$. Let $F_1, \dots, F_q \in \mathbb{R}[X]$ have degree at most k . If an interval I consists of*

$$r \geq \mathcal{P}_{k,q}(v, \varepsilon)$$

consecutive integers, then I has a partition into nonempty arithmetic progressions P satisfying

$$|P| \in \{v-1, v\}, \quad \text{osc}_P(F_i) \leq \varepsilon \quad (1 \leq i \leq q).$$

The number of cells is determined by the construction; it is not prescribed in advance.

[write] **6 October 2026; sources 15–20.** *Relation to the other partition results.* With one linear phase per degree this is source 13's prescribed-length theorem: in degree one the exponent recursion gives source 13's $1/(3q+2)$ (Corollary I.26.9). Source 19's Theorem I.34.3 proves the same transfer independently, with degree profiles and flexible lengths (note after Corollary I.32.4). The formal `lemma_5_9` (`Section05.lean`, line 261; exact companion `lemma_5_9_holds`, `Proofs05FullPartition.lean`, line 98) has the exponent $1/(2K^q)$ with $K = (k!)^2 2^{(k+1)^2}$; the improvement here is not formalized.

For $v = 1$, the analogous conclusion follows from the singleton partition, so the restriction $v \geq 2$ omits only a trivial case.

Proof. We induct on the degree. For degree one, constants can be discarded. Write $F_i(x) = \alpha_i x + \beta_i$ and put $Q = \lceil (v-1)/\varepsilon \rceil$. Place the $Q^q + 1$ points

$$0, \alpha, 2\alpha, \dots, Q^q \alpha \quad \text{in } \mathbb{R}/\mathbb{Z}^q, \quad \alpha = (\alpha_1, \dots, \alpha_q),$$

in Q^q half-open boxes of side $1/Q$. Two points lie in the same box, so their index difference gives an integer $1 \leq p \leq Q^q$ with $\|p\alpha_i\|_{\mathbb{R}/\mathbb{Z}} \leq 1/Q$ for every i .

Split I into residue chains modulo p . Each chain has length at least $\lfloor r/p \rfloor \geq g(v)$. By Lemma I.32.1, each chain can be subdivided into consecutive blocks of lengths $v-1$ or v . On any such block the i th phase admits a real lift with range length at most $(v-1)\|p\alpha_i\|_{\mathbb{R}/\mathbb{Z}} \leq \varepsilon$. This proves the base case, including the containing-arc conclusion.

Suppose the theorem has been proved through degree $j-1$. Write

$$F_i(x) = \alpha_i x^j + \text{terms of degree at most } j-1.$$

With U, T as in (I.32.3), choose $p \leq T$ using (I.32.1). Since $r \geq UT$, every residue chain modulo p has length at least U . Subdivide these chains into blocks of lengths in $[U, 2U-1]$.

On one block write $x = b + pt$, where $0 \leq t < n$ and $U \leq n < 2U$. Expansion gives

$$F_i(b + pt) = \alpha_i p^j t^j + G_i(t), \quad \deg G_i \leq j-1.$$

Choose a representative $c_i \in \mathbb{R}$ of $\alpha_i p^j \pmod{1}$ with

$$|c_i| \leq \frac{\varepsilon}{2^{j+1}U^j}.$$

The real numbers $c_i t^j$, $0 \leq t < n$, lie in an interval of length at most

$$(2U)^j |c_i| \leq \varepsilon/2.$$

Apply the induction hypothesis simultaneously to $G_1, \dots, G_q$ on $\{0, \dots, n-1\}$, with error $\varepsilon/2$. Its length condition holds because $n \geq U = \mathcal{P}_{j-1,q}(v, \varepsilon/2)$. Pull the resulting subprogressions back by $t \mapsto b + pt$. Their lengths are $v-1$ or v . On each one, the lower-degree part lies in an arc of length $\varepsilon/2$, and the leading part in an arc of length $\varepsilon/2$. Their sum lies in an arc of length ε . Performing this construction on all blocks gives a disjoint partition of the original interval. $\square$

The necessity of descending through the lower coefficients is real. A return $\|\alpha_i p^j\|_{\mathbb{R}/\mathbb{Z}} \ll 1$ controls the coefficient of t^j in $F_i(b + pt)$, but the translated lower coefficients contain the possibly large integer b . They do not become small merely because the leading coefficient does. The proof treats them as a new simultaneous lower-degree problem.

I.32.3 Power laws and their constants

Suppose the recurrence radii satisfy

$$\mathcal{R}_{j,q}(\eta) \leq C_j \eta^{-s_j} \quad (0 < \eta \leq 1), \quad C_j \geq 1, \quad s_j > 0. \quad (\text{I.32.4})$$

Dependence of C_j and s_j on q is allowed. Define

$$\begin{aligned} A_1 &= q + 2, & B_1 &= q, & D_1 &= 2^q, \\ A_j &= (1 + js_j)A_{j-1}, & B_j &= (1 + js_j)B_{j-1} + s_j, & & \\ D_j &= C_j 2^{(j+1)s_j + B_{j-1}(1+js_j)} D_{j-1}^{1+js_j}. \end{aligned} \quad (\text{I.32.5})$$

Proposition I.32.3 (A quantitative exponent region). *Under (I.32.4),*

$$\mathcal{P}_{k,q}(v, \varepsilon) \leq D_k v^{A_k} \varepsilon^{-B_k}.$$

Consequently, for every $a, b > 0$ with

$$A_k a + B_k b < 1, \tag{I.32.6}$$

all sufficiently large intervals of length r have the partition in Theorem I.32.2 with error r^{-b} , for every integer $2 \leq v \leq r^a$. Once the constants C_j are known, the sufficient threshold $r \geq D_k^{1/(1-A_k a - B_k b)}$ is explicit.

Proof. The inequalities $g(v) \leq v^2$ and $\lceil (v-1)/\varepsilon \rceil \leq 2v/\varepsilon$ prove the base case. Assuming the asserted bound in degree $j-1$, we have

$$U \leq D_{j-1} 2^{B_{j-1}} v^{A_{j-1}} \varepsilon^{-B_{j-1}}.$$

The defining recursion and (I.32.4) give

$$\mathcal{P}_{j,q}(v, \varepsilon) \leq C_j 2^{(j+1)s_j} U^{1+j s_j} \varepsilon^{-s_j}.$$

Substitution proves precisely (I.32.5). Finally, $D_k v^{A_k} r^{B_k b} \leq D_k r^{A_k a + B_k b} \leq r$ under the displayed threshold. $\square$

The exponents can also be written without recursion:

$$A_k = (q+2) \prod_{j=2}^k (1 + j s_j), \tag{I.32.7}$$

$$B_k = q \prod_{j=2}^k (1 + j s_j) + \sum_{i=2}^k s_i \prod_{j=i+1}^k (1 + j s_j). \tag{I.32.8}$$

An empty product is one. On the boundary of (I.32.6), one may retain the explicit prefactor

$$v \leq D_k^{-1/A_k} r^{(1-B_k b)/A_k}.$$

Thus the strict inequality in the simpler asymptotic statement is accounting for constants, not concealing an additional loss of a power.

The ratio in Gowers' Lemma 5.9 is $b = 2a$. Put $E_k = A_k + 2B_k$. Then

$$E_1 = 3q + 2, \quad E_j = (1 + j s_j) E_{j-1} + 2s_j. \tag{I.32.9}$$

Every $a < 1/E_k$ is therefore a valid prescribed-length exponent at that error ratio.

There is a useful distinction between prescribed adjacent lengths and lengths allowed to vary over a factor of two. If the desired final lengths are any integers in $[v, 2v-1]$, the degree-one threshold can instead be taken to be

$$v \left\lceil \frac{2v}{\varepsilon} \right\rceil^q.$$

Use the same Dirichlet step and the second subdivision rule in Lemma I.32.1. The rest of the induction is unchanged. In this variant $A_1 = q+1$ and $B_1 = q$; one may take $D_1 = 4^q$. The additional factor of v in the adjacent-length threshold is the cost of the elementary integer representation used to force those two particular lengths.

I.32.4 Consequences of modern simultaneous recurrence

Maynard's published theorem [72] gives, for fixed degree j , a constant $c_j > 0$ such that for any q real polynomials h_i of degree at most j with $h_i(0) = 0$, there is $1 \leq p < T$ with

$$\max_i \|h_i(p)\|_{\mathbb{R}/\mathbb{Z}} \ll_{j,q} T^{-c_j/q}.$$

Apply this to $h_i(X) = \alpha_i X^j$ and invert the error bound. After allowing for an initial threshold and integer rounding, it gives (I.32.4) with $s_j = q/c_j$. Equations (I.32.7)–(I.32.8) therefore give

$$A_k, B_k, E_k = O_k(q^k)$$

for each fixed degree. This conclusion uses an established analytic input; the contribution proved above is its exact partition transfer.

The preprint [73], Corollary 1.2, gives a more explicit degree dependence. For every fixed $\kappa > 10.5$, its bound permits

$$s_j = \kappa q j(j-1), \quad j \geq 2. \quad (\text{I.32.10})$$

The implied constants are uniform over the polynomial coefficients, but depend on j, q, κ ; they are not supplied here as numerical finite-size constants. We use this result as an imported preprint theorem and do not claim a new proof of it. Substitution yields

$$\begin{aligned} A_j &= (1 + \kappa q j^2(j-1))A_{j-1}, \\ B_j &= (1 + \kappa q j^2(j-1))B_{j-1} + \kappa q j(j-1). \end{aligned}$$

For example,

$$E_k = 3\kappa^{k-1}(k!)^2(k-1)!q^k + O_{k,\kappa}(q^{k-1}).$$

Corollary I.32.4 (Simultaneous quadratic partitions). *Let $q \geq 1$ and*

$$0 < a < \frac{1}{126q^2 + 129q + 2}.$$

For all sufficiently large r , every family of q real quadratic polynomials on an interval of r integers admits, for every integer $2 \leq v \leq r^a$, a partition into progressions of lengths $v-1$ or v on which all containing-arc oscillations are at most r^{-2a} . The threshold is uniform over the coefficients. The same conclusion holds for q modular quadratic polynomials, with diameters at most Nr^{-2a} , for every $N \geq r$.

Proof. For degree two, (I.32.10) gives $s_2 = 2\kappa q$, and hence

$$\begin{aligned} A_2 &= (q+2)(1+4\kappa q), \\ B_2 &= q(1+4\kappa q) + 2\kappa q, \\ E_2 &= 12\kappa q^2 + (3+12\kappa)q + 2. \end{aligned}$$

The strict hypothesis allows $\kappa > 10.5$ sufficiently close to 10.5 that $aE_2 < 1$. Apply Proposition I.32.3 with $b = 2a$, and then use the modular interpretation. $\square$

[write] **6 October 2026; sources 15–20.** *Source 13's Question 3 is answered* by this corollary and Corollary I.32.5 (for quadratics, exponents polynomial in the number of phases; for a shared leading coefficient, linear in q , from the source's own scalar Lemma 5.5), and independently by source 19's Theorem I.34.3 with Corollary I.34.5 (item I.1). Both rest on Maynard's published theorem; the numbers rest on Lau's preprint. *Two routes compared* (checked at the write): at the same recurrence constant this corollary's prescribed-length exponent is strictly better than source 19's, because source 19 squares the length only at the end (its Lemma I.34.4) while this proof pays v^2 once at the base: with $\kappa = 11$ the reciprocal exponents at $q = 1$ are 269 here and 314 in source 19 ($2(F_2 + G_2)$, source 19's Table I.34.2). Source 19's is more general (mixed degree profiles). Source 16's Theorem I.30.5 is a different problem (own steps, affine frequencies) with the same input.

Taking the convenient value $\kappa = 11$ gives the slightly weaker denominator $132q^2 + 135q + 2$ without a limiting choice of κ . In contrast, the degree-two value of Gowers' constant is $K = (2!)^2 2^9 = 2048$, so the displayed exponent in Lemma 5.9 is $1/(2 \cdot 2048^q)$ at the same ratio of diameter and length powers. The corollary improves this local dependence from exponential to polynomial in q . It does not, by itself, improve the paper's final bound for Szemerédi's theorem: the other multilinear partition and inverse-theorem losses must also be propagated before such a claim would be justified.

I.32.5 A special case requiring only scalar recurrence

The transfer can exploit a smaller number of independent leading frequencies even when the lower coefficients are unrestricted. Here is a concrete quadratic example.

Corollary I.32.5 (A shared leading coefficient). *Suppose a scalar recurrence estimate has the following form: for every $\alpha \in \mathbb{R}/\mathbb{Z}$ and integer $T \geq T_0$, where $T_0 \in \mathbb{Z}_{\geq 1}$, there is $1 \leq p \leq T$ with $\|\alpha p^2\|_{\mathbb{R}/\mathbb{Z}} \leq T^{-1/s}$, where $s > 0$. Let*

$$F_i(X) = h_i \alpha X^2 + b_i X + c_i \pmod{1}, \quad h_i \in \mathbb{Z}, \quad |h_i| \leq H,$$

where $H \geq 1$ and the lower coefficients are arbitrary. Put

$$U = g(v) \left\lceil \frac{2(v-1)}{\varepsilon} \right\rceil^q, \quad T = \max \left\{ T_0, \left\lceil \left(\frac{8HU^2}{\varepsilon} \right)^s \right\rceil \right\}.$$

Every integer interval of length at least UT has the partition of Theorem I.32.2. For fixed H, T_0, s, q , every exponent

$$a < \frac{1}{(1+2s)(3q+2)+2s}$$

works asymptotically with $v \leq r^a$ and error r^{-2a} .

Proof. The chosen T ensures $\|h_i \alpha p^2\|_{\mathbb{R}/\mathbb{Z}} \leq HT^{-1/s} \leq \varepsilon/(8U^2)$ for every i . The quadratic step of the transfer now applies. Its recurrence-radius exponent is s , independent of q , whereas the lower linear step still has q slopes. Equation (I.32.9) gives the stated denominator. $\square$

The ordinary degree-two recurrence exponent in the supplied version of Gowers' Lemma 5.5 is $1/16$. Taking that lemma with the explicit repaired threshold $T_0 = 2^{2^{320}}$ gives $s = 16$ and denominator $99q + 98$ in this special case. In particular the finite threshold in Corollary I.32.5 is $U \max\{2^{2^{320}}, \lceil (8HU^2/\varepsilon)^{16} \rceil\}$, with U as displayed there. Thus a linear reciprocal dependence on q is already available from scalar recurrence when the leading coefficient is shared; no modern simultaneous recurrence theorem is needed for that consequence. Bounded integer combinations

of s_0 leading frequencies similarly replace the recurrence dimension q by s_0 at the quadratic step, while retaining all q lower slopes.

For completeness, the modular scalar recurrence extends to an arbitrary real α at the same fixed integer T by approximating α with rational coefficients whose denominators tend to infinity. The finitely many possible values $1 \leq p \leq T$ include one occurring on an infinite subsequence, and continuity then gives the asserted torus recurrence.

I.33 A volume obstruction to polynomial partition exponents

Source 18, delivered Section 5.

The preceding sufficient exponents need not be optimal. We now prove an obstruction that applies to every universal theorem, regardless of the construction used to obtain its partition. Its proof is elementary and independent of all imported recurrence estimates.

I.33.1 Lifting a polynomial from a short arc

Lemma I.33.1 (Finite-difference rigidity). *Let $F \in \mathbb{R}[X]$ have degree at most k , let $F(0) = 0$, and let $n \geq k$. Suppose*

$$\|F(t)\|_{\mathbb{R}/\mathbb{Z}} \leq \varepsilon \quad (0 \leq t \leq n), \quad 0 < \varepsilon < 2^{-k-1}.$$

Let $y_t \in [-\varepsilon, \varepsilon]$ be the unique lift of $F(t) \pmod{1}$. Then

$$y_t = \sum_{j=1}^k b_j \binom{t}{j} \quad (0 \leq t \leq n), \quad b_j = \Delta^j y_0,$$

where $\Delta y_t = y_{t+1} - y_t$. Moreover $|b_j| \leq 2^j \varepsilon < 1/2$, and b_j is the canonical lift of the j th binomial coefficient of F modulo one.

Proof. The k th finite differences of the real lifts all lie in $[-2^k \varepsilon, 2^k \varepsilon] \subset (-1/2, 1/2)$. Modulo one, each equals the same k th finite difference of the polynomial F . Uniqueness of a representative in this interval therefore makes these real finite differences equal. The Newton formula for a finite sequence with constant k th differences gives the displayed polynomial representation. The formula

$$\Delta^j y_0 = \sum_{t=0}^j (-1)^{j-t} \binom{j}{t} y_t$$

gives $|b_j| \leq 2^j \varepsilon$. Reduction modulo one identifies b_j with the corresponding binomial coefficient of F , and the strict bound $|b_j| < 1/2$ makes the lift canonical. $\square$

Lemma I.33.2 (A binomial evaluation determinant). *For positive integers k, m , put $H = k(k+1)/2$. Then*

$$\det \left[\binom{im}{j} \right]_{1 \leq i, j \leq k} = m^H.$$

Proof. The polynomial $\binom{X}{j}$ has degree j , leading coefficient $1/j!$, and vanishes at zero for $j \geq 1$. The Vandermonde determinant at $0, m, 2m, \dots, km$ therefore gives

$$\det \left[\binom{im}{j} \right] = \frac{\prod_{i=1}^k im \prod_{1 \leq i < \ell \leq k} (\ell - i)m}{\prod_{j=1}^k j!} = m^{k(k+1)/2}.$$

The integer factors in the numerator cancel the factorials in the denominator. $\square$

I.33.2 The probability of a flat progression

Choose independent Haar-uniform coefficients $\alpha_{i,j} \in \mathbb{R}/\mathbb{Z}$, $1 \leq i \leq q$, $1 \leq j \leq k$, and interpret

$$F_i(X) = \sum_{j=1}^k \alpha_{i,j} X^j \pmod{1}$$

as random polynomial phases. Any real representatives define the same values at integers.

Theorem I.33.3 (A fixed-progression volume bound). *Let $k, q \geq 1$, $n \geq k$, and $0 < \varepsilon < 2^{-k-1}$. Put*

$$H = \frac{k(k+1)}{2}, \quad m = \left\lfloor \frac{n}{k} \right\rfloor.$$

For every positive integer p ,

$$\Pr \left[\text{osc}_{\{0,p,\dots,np\}}(F_i) \leq \varepsilon \text{ for all } i \right] \leq \left(\frac{(2\varepsilon)^k}{m^H} \right)^q. \quad (\text{I.33.1})$$

Proof. It suffices to treat a single polynomial. Express its values in the binomial basis:

$$F(pt) = \sum_{j=1}^k c_j \binom{t}{j} \pmod{1}.$$

The transformation from the monomial coefficients $(\alpha_1, \dots, \alpha_k)$ to $(c_1, \dots, c_k)$ is an integer triangular endomorphism of $\mathbb{R}/\mathbb{Z}^k$ with diagonal entries $j!p^j$. Its determinant is nonzero, so the endomorphism is surjective and sends Haar measure to Haar measure. In particular, the vector $c = (c_1, \dots, c_k)$ is uniform on $\mathbb{R}/\mathbb{Z}^k$.

On the event under consideration, $F(0) = 0$ and all values are within geodesic distance ε of zero. Apply Lemma I.33.1 to $t \mapsto F(pt)$. There are canonical representatives $b_j \in (-1/2, 1/2)$ of c_j such that the lifted values are

$$y_t = \sum_{j=1}^k b_j \binom{t}{j}, \quad |y_t| \leq \varepsilon.$$

Evaluate at $t = m, 2m, \dots, km$, all of which are at most n . By Lemma I.33.2, the linear map from b to these k values has determinant m^H . The conditions that all those values lie in $[-\varepsilon, \varepsilon]$ define a parallelootope of volume $(2\varepsilon)^k/m^H$. Intersecting it with the fundamental cube cannot increase the volume. There is no sum over different integer lifts: the rigidity lemma has already made the representative b unique. This proves the one-polynomial estimate, and independence of the q coefficient vectors gives (I.33.1). $\square$

The powers in this fixed-progression estimate are sharp up to constants depending on k . For one polynomial, the sufficient restrictions on its canonical binomial coefficients

$$|c_j| \leq \frac{\varepsilon}{2k \binom{n}{j}}, \quad 1 \leq j \leq k,$$

ensure that all lifted values lie in $[-\varepsilon/2, \varepsilon/2]$, and hence that their containing arc has length at most ε . The Haar measure of this rectangle is

$$\prod_{j=1}^k \frac{\varepsilon}{k \binom{n}{j}} \geq \frac{\prod_{j=1}^k j!}{k^k} \varepsilon^k n^{-H}.$$

On the other hand, $m \geq n/(2k)$ in (I.33.1). Thus the one-polynomial probability is

$$\Theta_k(\varepsilon^k n^{-k(k+1)/2})$$

throughout the stated small-arc range. Sharpness of this probability scale does not establish a matching universal partition theorem: one still has to cover every point by disjoint progressions, and the flatness events for different progressions are dependent.

I.33.3 Counterexamples to excessive universal exponents

Theorem I.33.4 (An explicit obstruction). *Let k, q, n, ε satisfy the hypotheses of Theorem I.33.3, let $r \geq n + 1$, and retain $H = k(k + 1)/2$. If*

$$\left\lfloor \frac{r-1}{n} \right\rfloor \left(\frac{(2\varepsilon)^k}{\lfloor n/k \rfloor^H} \right)^q < 1, \quad (\text{I.33.2})$$

then there are q polynomial phases of degree at most k for which $\{0, \dots, r-1\}$ has no partition into arithmetic progressions all of length at least $n+1$ and having all q oscillations at most ε . These counterexamples may also be realized by modular polynomials with a sufficiently large common modulus $N \geq r$.

Proof. Any arithmetic progression in $\{0, \dots, r-1\}$ containing zero can be enumerated as $0, p, 2p, \dots$ with positive step. If it has length at least $n+1$ and all q oscillations at most ε , its first $n+1$ points have the same property, and $p \leq \lfloor (r-1)/n \rfloor$. The union bound and Theorem I.33.3 show that the probability of any such progression is less than one. There is therefore a coefficient vector for which zero belongs to no such progression. The cell containing zero rules out the asserted partition.

For the modular assertion, there are only finitely many candidate steps. Each good-progression event is closed in the coefficient torus: it asks that a finite collection of continuously varying points fit into a closed arc of prescribed length. One may make the arc's starting point an auxiliary circle variable and project the resulting closed subset of a compact space. Consequently the nonempty complement just constructed is open. Approximate one of its points by rational coefficient vectors with a common denominator $N \geq r$ sufficiently large. The strict failure for every candidate step persists. Multiplication by N then gives polynomials over $\mathbb{Z}/N\mathbb{Z}$. $\square$

Using $\lfloor n/k \rfloor \geq n/(2k)$, a simpler sufficient condition for the same obstruction is

$$r \left[2^k (2k)^{k(k+1)/2} \varepsilon^k \right]^q < n^{qk(k+1)/2+1}. \quad (\text{I.33.3})$$

Corollary I.33.5 (A necessary exponent region). *Suppose $a, b, c, C > 0$, and suppose a universal assertion states that for every sufficiently large r and every q polynomial phases of degree at most k , an interval of r integers has a partition whose cells all have length at least cr^a and all q oscillations at most Cr^{-b} . Then necessarily*

$$\left(\frac{qk(k+1)}{2} + 1\right)a + kqb \leq 1. \quad (\text{I.33.4})$$

The same necessary condition holds if the universal assertion is restricted to modular polynomials and is uniform over $N \geq r$.

Proof. If (I.33.4) fails, choose an integer n comparable to r^a and just below the guaranteed minimum cell length, and put $\varepsilon = Cr^{-b}$. For large r the small-arc hypothesis holds, while the ratio of the left side of (I.33.3) to its right side tends to zero. The preceding theorem contradicts the universal assertion. If the proposed minimum length already exceeds r , the assertion is impossible directly. $\square$

[write] **6 October 2026; sources 15–20.** *Credit.* At $k = 1$ this is source 13’s frontier $(q+1)a + qb \leq 1$ (Theorem I.26.5, Corollary I.26.6) and source 05’s minimax order (Theorem I.18.5). Source 19’s Theorem I.34.7 is the same obstruction for arbitrary degree profiles, $(S+1)a + Db \leq 1$; at uniform degree k it is this one. Source 16’s Corollary I.31.2 is the own-step analogue.

For $k = 1$, the necessary region is $(q+1)a + qb \leq 1$, agreeing with the flexible linear partition region. For quadratics it is

$$(3q+1)a + 2qb \leq 1.$$

At the ratio $b = 2a$ this becomes $a \leq 1/(7q+1)$. More generally, that ratio gives

$$a \leq \frac{1}{qk(k+5)/2 + 1}.$$

These are universal obstructions, not constructions. In particular, the quadratic sufficient denominator $126q^2 + 129q + 2$ and the necessary denominator $7q + 1$ should not be presented as matching bounds. Closing their dimensional gap, while also distinguishing flexible cell lengths from prescribed adjacent lengths, is a concrete further research problem.

I.34 Simultaneous polynomial phase partitions

Source 19, delivered Section 6.

Lemma 5.9 of Gowers’ paper obtains simultaneous control of q polynomial phases by refining a partition once for each phase. Its exponent therefore contains K^{-q} , where $K = (k!)^2 2^{(k+1)^2}$ for degree at most k . The repository repairs the statement to prescribe a target cell length, with an emergent number of cells [1, 89]. Here we replace refinement by phase with refinement by degree. At a given degree, all leading coefficients are made small simultaneously. The resulting exponent has polynomial dependence on q for fixed k .

The substantial recurrence input is an established theorem of Maynard; the contribution proved here is its transfer to complete partitions, including the degree profile, rounding, and error budgets. We also prove a separate obstruction by random coefficients. That obstruction shows that the sufficient bounds below still leave a substantial gap.

I.34.1 The recurrence input and its precise status

Write $\|t\|_{\mathbb{R}/\mathbb{Z}}$ for distance to the nearest integer. We use the following published result.

Theorem I.34.1 (Maynard, Theorem 1.1 [72]). *For every positive integer j there is a constant $\kappa_j > 2$, independent of the number m of polynomials, and for every m a constant $U_{j,m} > 2$, with the following property. Let $f_1, \dots, f_m \in \mathbb{R}[X]$ have degree at most j and vanish at zero. If $0 < \varepsilon_i \leq 1/100$,*

$$\left(\prod_{i=1}^m \varepsilon_i \right)^{-1} \leq x^{1/\kappa_j}, \quad x > U_{j,m},$$

then some integer $1 \leq h < x$ satisfies $\|f_i(h)\|_{\mathbb{R}/\mathbb{Z}} \leq \varepsilon_i$ for every i .

The theorem is uniform in every polynomial coefficient. Maynard does not give numerical values for the constants needed here. In particular, his discussion that a certain numerical degree dependence would probably suffice is not a proved numerical choice. We keep the published constants symbolic in the principal result.

For degree one we improve the input to $\kappa_1 = 1$ using simultaneous Dirichlet approximation. More generally, all subsequent arguments work with any valid recurrence constants. Lau's preprint gives an optional numerical specialization, discussed separately below; it is not needed for the conclusion about polynomial dependence on the number of phases.

Lemma I.34.2 (Power-scale recurrence). *Set $\kappa_1 = 1$, and for $j \geq 2$ take the constants in Theorem I.34.1. If $e, s > 0$ and $\kappa_j m e < s$, then for all sufficiently large r , uniformly for $\eta_1, \dots, \eta_m \in \mathbb{R}$, there is an integer $1 \leq h \leq r^s$ with*

$$\|\eta_i h^j\|_{\mathbb{R}/\mathbb{Z}} \leq r^{-e} \quad (1 \leq i \leq m).$$

Proof. For $j \geq 2$, apply Theorem I.34.1 with $x = r^s$, $f_i(X) = \eta_i X^j$, and $\varepsilon_i = r^{-e}$. Its product condition follows from $\kappa_j m e < s$, and the other conditions hold once $r^e \geq 100$ and $r^s > U_{j,m}$. For $j = 1$, partition the unit m -cube into Q^m equal subcubes and apply the pigeonhole principle to the $Q^m + 1$ points $n(\eta_1, \dots, \eta_m) \bmod 1$, $0 \leq n \leq Q^m$. With $Q = \lceil r^e \rceil$, the difference of two points in one subcube gives $h \leq Q^m \leq 2^m r^{me}$ and errors at most $1/Q \leq r^{-e}$. The bound $h \leq r^s$ follows once $r \geq 2^{m/(s-me)}$. $\square$

I.34.2 A theorem sensitive to the degree profile

For a finite subset $S \subset \mathbb{R}/\mathbb{Z}$, let $\text{ad}(S)$ be the least length of a real interval whose projection contains S , allowing the upper bound 1. This circular arc diameter is monotone under restriction and satisfies

$$\text{ad}(S + T) \leq \text{ad}(S) + \text{ad}(T).$$

We apply it to polynomial values modulo one.

Let $q_1 \geq \dots \geq q_k \geq 0$ be integers. The interpretation is that at most q_j of the phases have degree at least j . Define

$$\begin{aligned} F_0 &= 1, & G_0 &= 0, \\ F_j &= (1 + j\kappa_j q_j)F_{j-1}, & G_j &= (1 + j\kappa_j q_j)G_{j-1} + \kappa_j q_j. \end{aligned} \tag{I.34.1}$$

Thus

$$F_k = \prod_{j=1}^k (1 + j\kappa_j q_j), \quad (\text{I.34.2})$$

$$G_k = \sum_{j=1}^k \kappa_j q_j \prod_{\ell=j+1}^k (1 + \ell\kappa_\ell q_\ell). \quad (\text{I.34.3})$$

Theorem I.34.3 (Complete polynomial phase partitions). *Let $a, b > 0$ satisfy*

$$F_k a + G_k b < 1. \quad (\text{I.34.4})$$

There is a threshold R , depending only on the displayed parameters and the recurrence constants, with the following property. For every integer $r \geq R$ and every family of real polynomials of degree at most k with the stated degree profile, the interval $[0, r) \cap \mathbb{Z}$ has a partition into nonempty ordinary arithmetic progressions P such that

$$|P| \geq r^a, \quad \text{ad}(\phi_i(P) \bmod 1) \leq r^{-b} \quad \text{for every phase } i \text{ and every cell } P.$$

The threshold is independent of the coefficients and constant terms. Different cells may have different common differences.

Proof. We induct on the maximal degree. For degree zero the phases are constant and the whole interval is a valid cell, since $a < 1$. At a positive degree k , discard an empty top degree if $q_k = 0$. Otherwise put $m = q_k$. The strict inequality (I.34.4) is

$$(1 + k\kappa_k m)(F_{k-1}a + G_{k-1}b) + \kappa_k mb < 1.$$

We can therefore choose $b' > b$, followed by τ and s , so that

$$F_{k-1}a + G_{k-1}b' < \tau, \quad \kappa_k m(k\tau + b') < s < 1 - \tau. \quad (\text{I.34.5})$$

These choices depend only on the theorem's parameters.

Write η_i for the leading coefficient of each degree- k phase, padding by zero coefficients if fewer than m phases have that degree. Apply Lemma I.34.2 with $e = k\tau + b'$. We obtain $1 \leq h \leq r^s$ and integers z_i with

$$\eta_i h^k = z_i + \delta_i, \quad |\delta_i| \leq r^{-(k\tau + b')}. \quad (\text{I.34.6})$$

Partition the original interval into residue chains modulo h . Every chain has at least $\lfloor r/h \rfloor \geq r^{1-s} - 1$ points. Put $L = \lceil r^\tau \rceil$. Since $s + \tau < 1$, each chain has at least L points once r is large enough. Split a chain into consecutive blocks of L points and merge any final remainder with the last full block. This produces subprogressions Q with lengths w satisfying

$$L \leq w < 2L \leq 4r^\tau. \quad (\text{I.34.7})$$

The blocks cover all points; no exceptional set is discarded.

On $Q = \{x, x+h, \dots, x+(w-1)h\}$, subtract the actual leading term $\eta_i h^k t^k$ from the polynomial $\phi_i(x+ht)$. The remainder $\psi_i(t)$ has degree at most $k-1$, and

$$\phi_i(x+ht) \equiv \delta_i t^k + \psi_i(t) \pmod{1} \quad (0 \leq t < w), \quad (\text{I.34.8})$$

because $z_i t^k$ is integer. For a phase of degree less than k , use $\delta_i = 0$ and simply reparameterize that phase. The coefficients of ψ_i can depend on the starting point x ; uniformity of the induction hypothesis permits this dependence.

At each lower degree $j < k$, the number of resulting phases of degree at least j is at most the original q_j . Indeed, affine reparameterization does not increase degree, and a degree- k phase has only been lowered to degree at most $k - 1$. Thus the lower profile is $q_1, \dots, q_{k-1}$, including the phases that have just lost their top degree.

Apply the induction hypothesis on $[0, w)$ with exponents a/τ and b'/τ . Their admissibility follows from (I.34.5). The resulting cells have at least $w^{a/\tau} \geq r^a$ points, and each lower phase has arc diameter at most $w^{-b'/\tau} \leq r^{-b'}$. On the same cell, the first term of (I.34.8) has arc diameter at most its real range, namely

$$|\delta_i|(w-1)^k \leq 4^k r^{-b'}.$$

Subadditivity gives a total bound of $(1 + 4^k)r^{-b'}$, which is at most r^{-b} once $r \geq (1 + 4^k)^{1/(b'-b)}$. Affine images of the child cells are ordinary arithmetic progressions, and all these images partition the original interval. $\square$

Two features of the proof matter. First, one cannot simply approximate all coefficients at the origin and assume that those approximations remain useful on every residue class: translation creates additional lower coefficients. Formula (I.34.8) handles those coefficients by a new, uniform lower-degree problem. Second, refinement runs through the degrees, rather than through the individual phases. When $q_j = q$, $F_k + G_k$ is a polynomial in q of degree k .

An explicit recursive threshold in the imported constants. The proof supplies a threshold rather than merely a coefficient-dependent “sufficiently large” condition. For a particular choice b', τ, s in (I.34.5), put $e = k\tau + b'$ and define

$$H_{k,m}(e, s) = \begin{cases} 2^{m/(s-me)}, & k = 1, \\ \max\{100^{1/e}, (U_{k,m} + 1)^{1/s}\}, & k \geq 2. \end{cases}$$

Starting with $R_0 = 2$, it suffices at degree k to choose an integer threshold at least

$$\max \left\{ 2, H_{k,m}(e, s), 4^{1/(1-s-\tau)}, (1 + 4^k)^{1/(b'-b)}, R_{k-1}(a/\tau, b'/\tau)^{1/\tau} \right\}. \quad (\text{I.34.9})$$

The lower-degree profile is understood in the last term. The third term guarantees enough points in every residue chain for (I.34.7), and the last term makes every child application legal. This expression is recursive and explicit in the imported $U_{j,m}$. It is not a numerical global threshold unless those external constants are themselves supplied numerically.

I.34.3 Prescribed target lengths and a numerical illustration

Lemma I.34.4 (The final rounding step). *If an arithmetic progression has $n \geq u^2$ points, with $u \geq 1$ integer, it can be partitioned into consecutive subprogressions of length u or $u + 1$.*

Proof. Write $n = tu + v$, $0 \leq v < u$. Then $t \geq u > v$. Use v blocks of size $u + 1$ and $t - v$ blocks of size u . $\square$

Corollary I.34.5 (A replacement for the target-length form of Lemma 5.9). *Under the assumptions of Theorem I.34.3, every integer target $v \geq 2$ with $(v - 1)^2 \leq r^a$ is attainable by a further partition into cells of length $v - 1$ or v , with the same phase-diameter bound. The target $v = 1$ is attained by singletons. In particular, for*

$$0 < \theta < \frac{1}{2(F_k + G_k)}, \quad (\text{I.34.10})$$

all sufficiently large r admit every target $v \leq r^\theta$, with diameters at most $r^{-2\theta}$.

Proof. Use Lemma I.34.4 with $u = v - 1$ inside each cell. Taking $a = b = 2\theta$ in Theorem I.34.3 gives the final assertion. Restriction to subsets does not increase diameter. $\square$

[write] **6 October 2026; sources 15–20.** *Second route.* Source 18’s Corollary I.32.4 proves the prescribed-length quadratic statement with the better exponent $1/(126q^2 + 129q + 2)$ at the same input (note there); this corollary’s degree profiles are more general. The formal `lemma_5_9` (`Section05.lean`, line 261; exact companion `lemma_5_9_holds`) has $\theta = 1/(2K^q)$; not formalized.

For comparison, the repository’s Lemma 5.9 has $\theta = 1/(2K^q)$ in this same ratio between length and diameter exponents [89]. For fixed k , the denominator in (I.34.10) is $O_k(q^k)$ rather than exponential in q . With the symbolic constants from Maynard this proves an improvement for all sufficiently large q ; it does not justify an uncomputed numerical comparison at every small q .

A numerical specialization from Lau’s preprint. Lau’s Theorem 1.1 permits recurrence constants $\kappa_j = c_2 j(j - 1)$ for $j \geq 2$, where $c_2 = 10.5 + 9m/(2m)^M + \varepsilon$, with the stated conditions on M and ε [73]. Choose $M = 5$ and $\varepsilon = 1/10$. These meet those conditions and give

$$c_2 \leq 10.5 + \frac{9}{32} + \frac{1}{10} = 10.88125 < 11.$$

We may consequently use the weaker, convenient constants $\kappa_j = 11j(j - 1)$, with an unspecified finite threshold. All numerical illustrations in the remainder of this subsection import that preprint statement. The published Maynard theorem alone supplies the symbolic version above.

For m nonconstant phases, of which s are quadratic and the remaining $m - s$ are linear, the degree profile is $(q_1, q_2) = (m, s)$. Since $\kappa_2 = 22$, formulas (I.34.2)–(I.34.3) give

$$F_2 = (m + 1)(44s + 1), \quad G_2 = m(44s + 1) + 22s. \quad (\text{I.34.11})$$

The case $s = 0$ reduces to the linear result. Table I.34.1 shows how the actual profile affects the estimate.

Family	q_1	q_2	F_2	G_2
Ten linear phases	10	0	11	10
Eight linear and two quadratic phases	10	2	979	934
Ten quadratic phases	10	10	4851	4630

Table I.34.1: Degree profiles with the optional input $\kappa_2 = 22$. The sufficient region is $F_2 a + G_2 b < 1$.

For q quadratic phases in particular,

$$F_2 + G_2 = 88q^2 + 68q + 1. \quad (\text{I.34.12})$$

Thus equal flexible length/error exponents satisfy $a = b < 1/(88q^2 + 68q + 1)$, and the target-length exponent in Corollary I.34.5 may be any number strictly less than the reciprocal of twice this denominator. For $k = 2$, Gowers' K is 2048. Table I.34.2 compares denominators in the target-length normalization. The table records exponent comparisons only; the respective finite thresholds differ, and the displayed new endpoints are not asserted to be attained.

q	$F_2 + G_2$	New denominator $2(F_2 + G_2)$	Old denominator $2K^q$
1	157	314	2^{12}
2	489	978	2^{23}
4	1681	3362	2^{45}
8	6177	12354	2^{89}
16	23617	47234	2^{177}

Table I.34.2: Quadratic target-length exponents, using Lau's preprint input. The new admissible exponents are strictly below the indicated reciprocals.

The theorem applies directly to modular phases by dividing polynomial coefficients by the modulus. It requires neither a prime modulus nor $r \leq N$ as an analytic statement on integer inputs. Using it inside a later part of the repository proof still requires checking every local width, threshold, and progression convention. These local improvements by themselves establish no new global bound for Szemerédi's theorem.

I.34.4 An obstruction from random polynomial coefficients

The sufficient frontier has degree- k polynomial dependence on the number of degree- k phases. A necessary condition has only linear dependence. The following elementary argument identifies that gap.

Lemma I.34.6 (A polynomial phase small on one progression). *Let $d \geq 1$, $L \geq d + 1$, and $0 < \varepsilon < 2^{-d-1}$. Choose $\alpha_1, \dots, \alpha_d$ independently and uniformly in $[0, 1)$, and put $P(X) = \sum_{j=1}^d \alpha_j X^j$. For every positive integer h ,*

$$\mathbb{P}\left(\|P(ht)\|_{\mathbb{R}/\mathbb{Z}} \leq \varepsilon \text{ for every } 0 \leq t < L\right) \leq (2\varepsilon)^d \left\lfloor \frac{L-1}{d} \right\rfloor^{-d(d+1)/2}. \quad (\text{I.34.13})$$

Proof. Write $F(t) = P(ht)$. On the event in question, choose the unique centered representative $y_t \in [-\varepsilon, \varepsilon]$ with $F(t) - y_t = m_t \in \mathbb{Z}$. Uniqueness holds because $\varepsilon < 1/2$, and $y_0 = 0$. Since F has degree at most d , every defined $(d+1)$ st difference satisfies

$$\Delta^{d+1}m_t = -\Delta^{d+1}y_t \in \mathbb{Z}, \quad |\Delta^{d+1}y_t| \leq 2^{d+1}\varepsilon < 1.$$

Thus $\Delta^{d+1}m_t = 0$. The finite integer sequence m_t is the restriction of the integer-valued polynomial

$$\sum_{j=0}^d \Delta^j m_0 \binom{t}{j}.$$

For $L = d + 1$ the same assertion follows directly by Newton interpolation; for larger L the vanishing differences propagate that interpolation. Consequently

$$y_t = \sum_{j=1}^d c_j \binom{t}{j}, \quad c_j = \Delta^j y_0, \quad |c_j| \leq 2^j \varepsilon < \frac{1}{2}. \quad (\text{I.34.14})$$

Now consider the torus vector

$$\beta_j = \Delta^j F(0) \pmod{1}, \quad 1 \leq j \leq d.$$

The linear map from $(\alpha_1, \dots, \alpha_d)$ to $(\beta_1, \dots, \beta_d)$ is an integer triangular matrix with diagonal entries $j!h^j$ and nonzero determinant. It is therefore a surjective endomorphism of $(\mathbb{R}/\mathbb{Z})^d$ and preserves normalized Haar measure. For example, this follows by integrating characters: every nonzero character pulls back to a nonzero character because the transpose matrix is injective over $\mathbb{Z}^d$. Hence β is uniform on the torus.

The determinant of this map need not be 1, and by itself does not justify a Euclidean volume estimate for a chosen lift. The required extra fact is (I.34.14): every c_j lies in $(-1/2, 1/2)$ and is congruent to β_j modulo one. Thus c is the unique representative of β in that fundamental cube. The probability of the event is bounded by the Euclidean volume of its permissible c -vectors in the cube, whose normalized volume is one.

Put $u = \lfloor (L-1)/d \rfloor$. The values $y_u, y_{2u}, \dots, y_{du}$ belong to $[-\varepsilon, \varepsilon]$. The matrix taking c to these values is $B_{ij} = \binom{iu}{j}$, $1 \leq i, j \leq d$. Because its columns are polynomials of degrees $1, \dots, d$, with zero constant terms and leading coefficients $1/j!$, the Vandermonde determinant gives

$$\begin{aligned} \det B &= \frac{\prod_{i=1}^d (iu) \prod_{1 \leq i < \ell \leq d} (\ell u - iu)}{\prod_{j=1}^d j!} \\ &= u^{d(d+1)/2}. \end{aligned}$$

The inverse image of the cube $[-\varepsilon, \varepsilon]^d$ under B has volume $(2\varepsilon)^d / \det B$. It contains all permissible c -vectors, proving (I.34.13). $\square$

Theorem I.34.7 (A necessary exponent inequality). *Fix positive degrees $d_1, \dots, d_q$, and set*

$$D = \sum_{i=1}^q d_i, \quad S = \sum_{i=1}^q \frac{d_i(d_i + 1)}{2}.$$

If $a, b > 0$ satisfy

$$(S+1)a + Db > 1, \tag{I.34.15}$$

then for all sufficiently large r there are polynomials of these degrees for which no partition of $[0, r) \cap \mathbb{Z}$ into arithmetic progressions can have both $|P| \geq r^a$ and $\text{ad}(\phi_i(P) \pmod{1}) \leq r^{-b}$ for every i, P . The phases can be chosen with rational coefficients, and can also be realized modulo a sufficiently large prime.

Proof. If $a > 1$, the length requirement is already impossible. Assume $0 < a \leq 1$, put $L = \lceil r^a \rceil$, and set $\varepsilon = r^{-b}$. Independently choose all nonconstant coefficients of each polynomial uniformly modulo one, with constant term zero. For large r the hypotheses of Lemma I.34.6 hold for every d_i , and $\lfloor (L-1)/d_i \rfloor \geq L/(4d_i)$.

Consider one progression $\{0, h, \dots, (L-1)h\}$ contained in $[0, r)$. If its phase arc diameter is at most ε , then every phase value has circular distance at most ε from the value at zero, which is zero. Lemma I.34.6, independently for the q phases, bounds the probability of this event by

$$K\varepsilon^D L^{-S}, \quad K = \prod_{i=1}^q \left[2^{d_i} (4d_i)^{d_i(d_i+1)/2} \right].$$

There are at most $(r - 1)/(L - 1) \leq 2r/L$ admissible positive h . A union bound therefore bounds the probability that any such progression has all the required small diameters by

$$2Kr\varepsilon^D L^{-S-1}. \quad (\text{I.34.16})$$

This tends to zero under (I.34.15). Thus some tuple of polynomials has no such progression through zero.

Every complete partition into progressions of lengths at least L has a cell containing zero. Orient its common difference positively. Since the ambient interval contains no negative points, this cell starts at zero; its first L points give one of the forbidden progressions. This is a contradiction.

For fixed r, L , only finitely many progressions through zero must be excluded. The arc diameter of a fixed finite set of polynomial values is continuous in the coefficients. Failure of a weak diameter condition is therefore open, and the tuple just obtained has a neighborhood consisting of counterexamples. This neighborhood contains rational coefficient tuples of the prescribed degrees. A common denominator realizes them as normalized modular polynomials. Alternatively, take a sufficiently large prime $p \geq \max\{r, d_1 + 1, \dots, d_q + 1\}$ and approximate every coefficient by a multiple of $1/p$ within that neighborhood. The resulting phases are normalized polynomials modulo p and retain the obstruction. $\square$

[write] **6 October 2026; sources 15–20.** *Credit.* At uniform degree k this is source 18’s Corollary I.33.5; at $k = 1$ it is source 13’s frontier (Theorem I.26.5). The general degree profile is new here.

[write] **6 October 2026; sources 34–39.** For the prescribed lengths L or $L + 1$ (every target), source 35 adds the numerical-semigroup conductor and obtains $(W + 2)a + \sum_i d_i b_i \leq 1$, $W = S$ (quadratics at $b = 2a$: $1/(7\nu + 2)$), by explicit phases forcing $D \mid h$ (its Corollary I.61.6 and the note there); the obstruction here, for flexible lengths, is unchanged.

In degree-profile notation,

$$D = \sum_{j=1}^k q_j, \quad S = \sum_{j=1}^k j q_j.$$

Thus a necessary condition for any uniform theorem of this kind is

$$\left(1 + \sum_{j=1}^k j q_j\right) a + \left(\sum_{j=1}^k q_j\right) b \leq 1. \quad (\text{I.34.17})$$

For q linear phases this is $(q + 1)a + qb \leq 1$, matching the strict sufficient region supplied by Dirichlet approximation. For q phases of degree k , it becomes

$$\left(1 + \frac{qk(k + 1)}{2}\right) a + qkb \leq 1.$$

We establish no matching nonlinear partition theorem. In particular, the exponent denominator $O_k(q^k)$ from the recurrence transfer remains larger than the $O_k(q)$ obstruction scale.

I.34.5 Finite diagnostics and a concrete obstruction

The accompanying exact-arithmetic script checks the determinant identity in Lemma I.34.6 for degrees 1 through 8 and five spacings, for a total of 40 cases. It also exhaustively checks all 101^2 quadratic coefficient pairs modulo 101, at each of the two spacings 1 and 7, on five-point

progressions with $\varepsilon = 1/9 < 1/8$. At both spacings it finds 61 centered-small events and 45 events with arc diameter at most $1/9$; every centered event has the asserted Newton lift and vanishing third differences. The finite coefficient-to-Newton map is bijective in these tests because 101 is coprime to its determinant. These checks are diagnostics of the proof's bookkeeping; an atomic finite grid is not substituted for the continuous Haar-volume argument.

For a concrete example, consider

$$\phi(x) = \frac{690x + 294x^2}{1009} \pmod{1}, \quad x \in [0, 80) \cap \mathbb{Z}.$$

The 11 candidate eight-point progressions starting at zero have common differences $1, \dots, 11$. Exact enumeration gives minimum circular arc diameter $523/1009$, which exceeds $1/9$. Therefore no partition of $[0, 80)$ into arithmetic progressions of at least eight points has phase arc diameter at most $1/9$ on every cell. This finite witness illustrates the endpoint-cell mechanism used in Theorem I.34.7; the asymptotic obstruction itself is proved by (I.34.16).

I.34.6 Questions left by the partition results

1. Can the dependence on q in the degree- k sufficient frontier be reduced from order q^k to order q , or even by one power of q ? Theorem I.34.7 leaves room for such a gain.
2. Is the boundary (I.34.17) achievable for nonlinear phases? It is sharp for linear phases, but the present work provides neither sharpness nor a counterexample in higher degree.
3. Can simultaneous recurrence be made sufficiently explicit in both degree and dimension to support numerical thresholds throughout the repository's density-increment argument?
4. Can prescribed target lengths be obtained without the squaring loss in Lemma I.34.4, while covering every point?
5. Can bounded-height rational dependencies among leading coefficients replace q_j by a smaller effective dimension in the recurrence step?

I.35 The partition problem and its arithmetic parameter

Source 23, delivered Section 2.

Write $I_R = \{0, 1, \dots, R-1\}$ and $\mathbb{T} = \mathbb{R}/\mathbb{Z}$. A progression in I_R is an ordinary integer progression

$$P = \{a, a + d, \dots, a + (\ell - 1)d\} \subset I_R, \quad d \geq 1.$$

Its length is its cardinality ℓ . All partitions below are disjoint and cover their indicated domain exactly, unless an exceptional set is explicitly allowed. Distinct cells may have different common differences.

For a finite subset $X \subset \mathbb{T}$, let $\text{adiam}(X)$ be the length of the shortest closed circular arc containing X . For example, $\text{adiam}\{0, 1/3, 2/3\} = 2/3$, although the largest pairwise circular distance is $1/3$. This is the normalized version of Gowers's definition immediately before Lemma 2.3: the smallest s for which a set in $\mathbb{Z}/N\mathbb{Z}$ lies in $\{a, a + 1, \dots, a + s\}$ [1]. This distinction will matter in the lower bound.

Fix $q \geq 1$ and affine phases

$$\phi_i(n) = \alpha_i n + \beta_i \pmod{1}, \quad 1 \leq i \leq q.$$

Let $\mathcal{S} \subset \mathbb{N}$ be a finite set of allowed lengths, and write

$$L = \min \mathcal{S} \geq 2, \quad H = \max \mathcal{S}.$$

The additive semigroup generated by these lengths is

$$\langle \mathcal{S} \rangle = \left\{ \sum_{s \in \mathcal{S}} n_s s : n_s \in \mathbb{Z}_{\geq 0} \right\}.$$

If $\gcd(\mathcal{S}) = 1$, its *conductor* $C = C(\mathcal{S})$ is the least integer such that every integer at least C belongs to this semigroup. Consequently $C - 1 \notin \langle \mathcal{S} \rangle$. Our assumption $L \geq 2$ implies $C \geq L$.

For completeness, finiteness of the conductor needs no deep theorem. The residues modulo L attained by the semigroup form a subgroup of $\mathbb{Z}/L\mathbb{Z}$: any additive subsemigroup of a finite group is a subgroup, since repeated addition supplies inverses. The gcd assumption makes this subgroup the whole group. Choose a nonnegative representable integer in each residue class. Every sufficiently large integer is one of these representatives plus a nonnegative multiple of L .

Definition I.35.1 (Universal eventual threshold). For $\epsilon = (\epsilon_1, \dots, \epsilon_q)$ with $0 < \epsilon_i < 1/2$, define $T_{\mathcal{S}}(\epsilon)$ to be the least positive integer R_0 such that for every $R \geq R_0$ and every choice of the q affine phases, I_R has a partition into progressions P satisfying

$$|P| \in \mathcal{S}, \quad \text{diam}(\phi_i(P)) \leq \epsilon_i \quad (1 \leq i \leq q).$$

The word “eventual” makes the quantifiers explicit. Particular values of R below the threshold can still be easy; for example, a useful divisibility relation may make an exact tiling possible. The threshold requires one guarantee valid at *every* later interval length. The upper bound below will prove that it is finite. If $\gcd(\mathcal{S}) > 1$, such an eventual threshold cannot be finite: infinitely many total cardinalities are not sums of allowed lengths, even for constant phases.

I.36 An explicit theorem with matching upper and lower scales

Source 23, delivered Section 3.

Theorem I.36.1 (Semigroup bounds for simultaneous phase partitions). *Let $\mathcal{S}, L, H, C, q$ be as above, with $\gcd(\mathcal{S}) = 1$. For $0 < \epsilon_i < 1/2$, put*

$$B = \left\lceil \frac{(H-1)^q}{\epsilon_1 \cdots \epsilon_q} \right\rceil - 1, \quad M_i = \left\lceil \frac{L-1}{\epsilon_i} \right\rceil - 1.$$

Then

$$\boxed{C \prod_{i=1}^q M_i \leq T_{\mathcal{S}}(\epsilon) \leq CB.} \tag{I.36.1}$$

The upper construction uses one common difference for every cell. The lower bound holds even when the difference is chosen separately for each cell. The lower examples can also be realized by affine maps on a prime cyclic group, restricted to a proper interval.

[write] **6 October 2026; sources 21–27.** *Credit.* For $\mathcal{S} = \{v-1, v\}$ the conductor is $(v-1)(v-2)$, so the upper bound with one common difference is source 13's Theorem I.26.8 (the source reproduces it with attribution, Section I.38); the matching lower bound, valid for cell-dependent differences, is new in this report.

The proof separates approximation of the phases from exact tiling of the resulting residue chains. The upper construction is the familiar Dirichlet argument, with the conductor recording the exact tiling requirement. The lower construction makes that requirement unavoidable.

I.36.1 Simultaneous approximation with unequal tolerances

We first record a slightly sharper finite form of the usual rectangular pigeonhole estimate. Its only additional ingredient is translation invariance of volume on the finite-dimensional torus.

Lemma I.36.2 (Torus packing estimate). *For $\alpha_i \in \mathbb{T}$ and $0 < \Delta_i < 1/2$, there is an integer*

$$1 \leq d \leq \left\lceil (\Delta_1 \cdots \Delta_q)^{-1} \right\rceil - 1 \quad \text{with} \quad \|d\alpha_i\|_{\mathbb{T}} \leq \Delta_i \quad \text{for all } i.$$

Proof. Let $K = \lceil (\prod_i \Delta_i)^{-1} \rceil$ and place a closed rectangular box of coordinate side lengths Δ_i at each of the K orbit points $n\alpha$, $0 \leq n < K$. Intersection of two boxes implies all the required coordinate inequalities for their index difference. If there is no such pair, the finitely many compact boxes are disjoint with positive separation. Enlarge every side by a sufficiently small positive amount so that they remain disjoint and all side lengths stay below one. Their total volume is then strictly greater than $K \prod_i \Delta_i \geq 1$, impossible in a torus of volume one. Thus two boxes intersect, and their difference is at most $K - 1$. $\square$

The weak inequality in the conclusion handles the exact-volume boundary. For a completely discrete construction one can instead use the following slightly less sharp integer bound.

Lemma I.36.3 (Rectangular pigeonhole principle). *Given $\alpha_1, \dots, \alpha_q \in \mathbb{T}$ and positive integers $Q_1, \dots, Q_q$, there is an integer*

$$1 \leq d \leq Q_1 \cdots Q_q \quad \text{such that} \quad \|d\alpha_i\|_{\mathbb{T}} < Q_i^{-1} \quad \text{for every } i.$$

Proof. Partition $[0, 1)^q$ into $Q_1 \cdots Q_q$ half-open rectangular boxes, with side lengths Q_i^{-1} in coordinate i . Among the $1 + Q_1 \cdots Q_q$ points $n(\alpha_1, \dots, \alpha_q)$, $0 \leq n \leq Q_1 \cdots Q_q$, two lie in the same box. Their positive index difference is the required d . $\square$

Proof of the upper bound in Theorem I.36.1. Apply Lemma I.36.2 with $\Delta_i = \varepsilon_i/(H-1)$ to obtain $d \leq B$. If $R \geq CB$, each residue class modulo d meets I_R in a chain with at least $\lfloor R/d \rfloor \geq C$ elements. Every chain length is therefore a sum of elements of $\mathcal{S}$. Cut the chain consecutively into blocks having those lengths.

Choose the centered representative θ_i of $d\alpha_i$. On any resulting cell of length $\ell \leq H$, its phase has a real lift with successive differences θ_i , and hence span

$$(\ell - 1)|\theta_i| \leq \frac{\ell - 1}{H - 1} \varepsilon_i \leq \varepsilon_i.$$

Reduction modulo one gives the required containing arc. The residue chains partition I_R , and cutting each chain preserves exact coverage. $\square$

I.36.2 Flatness forces a small step

Lemma I.36.4 (Rigidity inside a short arc). *Suppose $\ell \geq 2$ and $c, c + \theta, \dots, c + (\ell - 1)\theta \in \mathbb{T}$ lie in an arc of length $\varepsilon < 1/2$. Then*

$$\|\theta\|_{\mathbb{T}} \leq \frac{\varepsilon}{\ell - 1}.$$

Proof. Lift every point to one real interval of length ε containing the arc. Each consecutive lifted difference belongs to $[-\varepsilon, \varepsilon]$ and is congruent to θ modulo one. Two such differences differ by an integer of absolute value at most $2\varepsilon < 1$, so they are equal. Call their common value t . The first and last lifts differ by $(\ell - 1)t$, proving $(\ell - 1)|t| \leq \varepsilon$. Since $|t| < 1/2$, it is the centered representative of θ . $\square$

Remark I.36.5 (An alternative diameter convention). If diameter means largest pairwise circular distance, the lemma remains true for $\varepsilon < 1/3$. Indeed, a set of that pairwise diameter lies in an arc of the same length: after placing one point at zero, the largest positive and negative centered representatives cannot have sum greater than ε , since their sum is at most $2\varepsilon < 1 - \varepsilon$. At $\varepsilon = 1/3$, the three equally spaced points give a counterexample. We use arc diameter throughout, so this extra restriction is unnecessary.

I.36.3 A mixed-radix obstruction

Proof of the lower bound in Theorem I.36.1. The integers M_i are at least two and satisfy

$$\frac{1}{M_i} > \frac{\varepsilon_i}{L - 1}.$$

Set $D_0 = 1$, $D_i = M_1 \cdots M_i$, and $D = D_q$. Take

$$\alpha_i = \frac{1}{D_i} \pmod{1}, \quad R = DC - 1. \tag{I.36.2}$$

Suppose an allowed flat progression has common difference d . If $D \nmid d$, choose the least i with $D_i \nmid d$. Then $D_{i-1} \mid d$, and the nonzero residue of d/D_{i-1} modulo M_i gives

$$\left\| \frac{d}{D_i} \right\|_{\mathbb{T}} \geq \frac{1}{M_i} > \frac{\varepsilon_i}{L - 1}.$$

This contradicts Lemma I.36.4, because the cell has at least L points. Thus every allowed cell has $D \mid d$ and lies in one residue class modulo D .

The last residue class, $D - 1$ modulo D , has exactly $C - 1$ points in I_R ; the other classes have C points. A partition would express $C - 1$ as a sum of allowed cell lengths, contradicting the defining gap of the conductor. We have exhibited a bad interval of length $DC - 1$, so the eventual threshold is at least DC . $\square$

The use of several phases is essential to the full product in the lower bound. The first phase checks divisibility by M_1 ; the next checks the next digit, and so on. This forces the same residue obstruction on every cell without assuming a common difference in the proposed partition.

I.36.4 Prime cyclic realization

Proposition I.36.6 (Robustness of the obstruction). *Let R, D_i, M_i be those in (I.36.2), and put*

$$\gamma = \min_i \left(\frac{1}{M_i} - \frac{\varepsilon_i}{L-1} \right) > 0.$$

For any prime p with

$$p \geq R, \quad p > \frac{R-1}{2(L-1)\gamma}, \quad (\text{I.36.3})$$

choose integers a_i nearest to p/D_i . The affine maps $x \mapsto a_i x$ on $\mathbb{Z}/p\mathbb{Z}$, restricted to I_R , admit no partition as in Theorem I.36.1, with normalized diameters ε_i .

Proof. An ordinary progression in I_R of length at least L has $d \leq (R-1)/(L-1)$. Also $|a_i/p - 1/D_i| \leq 1/(2p)$. If $D \nmid d$, its first failing digit therefore satisfies

$$\left\| \frac{da_i}{p} \right\|_{\mathbb{T}} \geq \left\| \frac{d}{D_i} \right\|_{\mathbb{T}} - \frac{d}{2p} > \frac{1}{M_i} - \gamma \geq \frac{\varepsilon_i}{L-1}.$$

Rigidity again forces $D \mid d$, and the same cardinality contradiction finishes the proof. Arbitrarily large primes exist, so the required prime can always be chosen. $\square$

This statement concerns an ordinary interval inside a prime cyclic group. It does not identify the interval length with the ambient prime. That is exactly the $r \leq N$ interval setting of Lemma 5.9.

I.37 Sharp asymptotics and the permitted range of cell sizes

Source 23, delivered Section 4.

I.37.1 Two consecutive lengths

The semigroup generated by L and $L+1$ has conductor $L(L-1)$. One elementary proof is to write $n = kL + r$, $0 \leq r < L$: if $k \geq L-1$, then

$$n = (k-r)L + r(L+1).$$

On the other hand, $L(L-1) - 1$ is not representable, as is also immediate from the residue $L-1$ modulo L .

Corollary I.37.1 (Sharp threshold for nearly equal cells). *For $\mathcal{S} = \{L, L+1\}$,*

$$L(L-1) \prod_{i=1}^q \left(\left\lceil \frac{L-1}{\varepsilon_i} \right\rceil - 1 \right) \leq T_{\{L, L+1\}}(\boldsymbol{\varepsilon}), \quad (\text{I.37.1})$$

$$T_{\{L, L+1\}}(\boldsymbol{\varepsilon}) \leq L(L-1) \left(\left\lceil \frac{L^q}{\varepsilon_1 \cdots \varepsilon_q} \right\rceil - 1 \right). \quad (\text{I.37.2})$$

For fixed q , as $L \rightarrow \infty$,

$$\boxed{T_{\{L, L+1\}}(\boldsymbol{\varepsilon}) = \left(1 + O_q(L^{-1})\right) \frac{L^{q+2}}{\varepsilon_1 \cdots \varepsilon_q}.} \quad (\text{I.37.3})$$

The error estimate is uniform over all $0 < \varepsilon_i < 1/2$, including tolerances tending to zero with L .

Proof. The finite inequalities are Theorem I.36.1. Uniformly in the indicated tolerances,

$$\left\lceil \frac{L-1}{\varepsilon_i} \right\rceil - 1 = \frac{L}{\varepsilon_i}(1 + O(L^{-1})), \quad \left\lceil \frac{L^q}{\prod_i \varepsilon_i} \right\rceil - 1 = \frac{L^q}{\prod_i \varepsilon_i}(1 + O(L^{-1})).$$

The conductor is $L^2(1 + O(L^{-1}))$. Multiplication and squeezing prove the assertion. $\square$

This gives the leading constant one, not just matching exponents. It also settles the question whether a clever choice of different progression differences can remove the quadratic tiling cost in the worst case. The lower construction allows all such choices and still forces them into residue classes of a forbidden cardinality.

I.37.2 Arbitrary narrow sets of lengths

Corollary I.37.2 (Narrow length sets). *For fixed q , let $\mathcal{S}$ vary with $L = \min \mathcal{S} \rightarrow \infty$, $\gcd(\mathcal{S}) = 1$, and $H/L \rightarrow 1$. Then, uniformly for $0 < \varepsilon_i < 1/2$,*

$$T_{\mathcal{S}}(\varepsilon) \sim \frac{C(\mathcal{S})L^q}{\varepsilon_1 \cdots \varepsilon_q}.$$

Proof. Divide (I.36.1) by $CL^q/(\prod_i \varepsilon_i)$. A lower bound for the ratio is

$$\left(1 - \frac{3}{2L}\right)^q,$$

and an upper bound is

$$\left(1 + \frac{H - L - 1}{L}\right)^q.$$

Both tend to one. These displayed bounds also give a version with growing q : it is enough that $q(H - L + 1)/L \rightarrow 0$. $\square$

The conductor can depend strongly on which lengths are allowed, even when their minimum and maximum are close. For example, for the two coprime lengths L and $L + s$, their conductor is $(L - 1)(L + s - 1)$; if every intermediate length is allowed, the next proposition gives a conductor of order L^2/s . Exact tiling therefore detects more than just the width of the allowed range.

I.37.3 An interval of permitted lengths

The following classical conductor computation is a consecutive-integer case of the Frobenius problem for arithmetic progressions [97]. We give the full argument because it directly determines the partition scale.

Proposition I.37.3 (Conductor of a consecutive interval). *For $1 \leq w \leq L - 1$,*

$$C(\{L, L + 1, \dots, L + w\}) = L \left\lceil \frac{L - 1}{w} \right\rceil. \quad (\text{I.37.4})$$

Proof. Sums of exactly k permitted lengths fill every integer in $[kL, k(L + w)]$: distribute any excess between zero and kw among k slots, each with capacity w . The k th and $(k + 1)$ st intervals touch or overlap exactly when $kw \geq L - 1$. Put $k_0 = \lceil (L - 1)/w \rceil$. All intervals from k_0 onward join to cover $[k_0 L, \infty)$. Meanwhile $k_0 L - 1$ lies above the end of the $(k_0 - 1)$ st interval and below the start of the k_0 th. Earlier intervals have still smaller endpoints. It is the last gap. $\square$

Corollary I.37.4 (The cost of reducing length variation). *For fixed q and $1 \leq w \leq L - 1$,*

$$T_{\{L, \dots, L+w\}}(\varepsilon) \asymp_q \frac{L^{q+2}}{w \varepsilon_1 \cdots \varepsilon_q}. \quad (\text{I.37.5})$$

If $w = o(L)$ as $L \rightarrow \infty$, then the sharper statement is

$$T_{\{L, \dots, L+w\}}(\varepsilon) \sim \frac{L^{q+2}}{w \varepsilon_1 \cdots \varepsilon_q}. \quad (\text{I.37.6})$$

Proof. Use (I.37.4) in Theorem I.36.1. Here $L \leq H < 2L$, and the conductor is comparable to L^2/w . If $w = o(L)$, its ceiling has relative error $O(w/L)$ and $H/L \rightarrow 1$, so Corollary I.37.2 gives the leading constant. $\square$

For equal tolerances ε , a logarithmic parametrization $L = X^a$, $\varepsilon = X^{-b}$, and $w = X^c$ gives the scale

$$T \asymp X^{(q+2)a+qb-c}, \quad 0 \leq c \leq a,$$

up to integer rounding and fixed factors. Thus the sharp scaling boundary is $(q+2)a + qb - c = 1$. Consecutive lengths have $c = 0$; a range of lengths comparable with L has $c = a$. The latter recovers the boundary $(q+1)a + qb = 1$ for flexible cell sizes. This is a statement about the threshold as its input parameters vary; it does not assert failure at every individual length below that threshold.

Table I.37.1: Dependence on the allowed lengths, for fixed q . The symbol $\sim$ includes leading constant one; $\asymp_q$ allows constants depending on q . All tolerances lie in $(0, 1/2)$.

Allowed lengths	Conductor	Eventual threshold
$\{L, L+1\}$	$L(L-1)$	$\sim L^{q+2} / \prod_i \varepsilon_i$
$\{L, \dots, L+w\}, w = o(L)$	$L \lceil (L-1)/w \rceil$	$\sim L^{q+2} / (w \prod_i \varepsilon_i)$
$\{L, \dots, 2L-1\}$	L	$\asymp_q L^{q+1} / \prod_i \varepsilon_i$
$\gcd(\mathcal{S}) > 1$	No finite conductor	No finite exact threshold

I.38 Consequences for Lemma 5.9 and exact finite witnesses

Source 23, delivered Section 5.

I.38.1 Which previously sufficient exponent is now optimal?

In Gowers's degree-one specialization, $K = 16$; the simultaneous lemma has target lengths $v-1, v$ with $v \leq R^{1/(2 \cdot 16^q)}$ and diameter $R^{-1/16^q}$. The repository has repaired the printed exact cell-count wording to this prescribed-length interface [107].

The earlier report *Sharp Local Refinements in Gowers' Szemerédi Framework*, Theorem 3.8, already gave the upper bound

$$R \geq (v-1)(v-2) \prod_i \left\lceil \frac{v-1}{\varepsilon_i} \right\rceil,$$

and its Corollary 3.9 gave the exponent $a_q = 1/(3q + 2)$ with phase diameter R^{-2a_q} [111]. This upper result is not claimed as new here. Its Question 11.1 asked whether the extra tiling cost was intrinsic. Theorem I.36.1 answers that question in the affirmative.

It is important to compare equal normalizations. If target length is R^a and diameter is R^{-b} , the prescribed-length boundary is

$$(q + 2)a + qb = 1. \quad (\text{I.38.1})$$

Under Gowers's relation $b = 2a$, this gives $a = 1/(3q + 2)$. Under the different relation $b = a$, it gives $a = 1/(2q + 2)$. Those are two slices of the same result, not competing improvements.

Proposition I.38.1 (Optimality for the all-targets interface). *Fix $q \geq 1$, $a > 0$, $b \geq 0$ and $0 < c < 1/2$. Suppose $(q + 2)a + qb > 1$. There are arbitrarily large R , q affine phases, and an integer $v \leq R^a$ such that no partition into cells of lengths $v - 1, v$ has all phase diameters at most cR^{-b} . The differences may be chosen separately on different cells.*

Proof. Choose $t > 0$ with

$$a(qt + 2) > 1, \quad t(1 - qb) < 1 + 2b.$$

Such t exists. If $qb \geq 1$, the second inequality is automatic. If $qb < 1$, the available interval for t is nonempty precisely because $a(q + 2) > 1 - qb$.

Let $L \rightarrow \infty$, $M = \lfloor L^t \rfloor$, $D = M^q$, and $R = D(L^2 - L - 1)$. Set $v = L + 1$ and $\alpha_i = M^{-i}$. Since $R \sim L^{qt+2}$, the first inequality gives $v \leq R^a$ eventually. Moreover,

$$\frac{McR^{-b}}{L - 1} \sim cL^{t(1-qb)-1-2b} \longrightarrow 0.$$

Thus $1/M > cR^{-b}/(L - 1)$ eventually. The same mixed-radix argument forces every admissible step to be divisible by D . Each residue has $L^2 - L - 1$ elements, the last gap of $\langle L, L + 1 \rangle$. $\square$

[write] **6 October 2026; sources 21–27.** *Answers source 13's Question 1* (item I.12). The prescribed-length region $(q + 2)a + qb \leq 1$ of source 13's Corollary I.26.9 cannot be enlarged, even with cell-dependent differences, so at Gowers's $b = 2a$ its exponent $1/(3q + 2)$ is optimal; the flexible frontier $(q + 1)a + qb = 1$ needs an exceptional set or flexible lengths. The existence of t in the proof was re-derived at the write: for $qb < 1$ the two conditions say $(1/a - 2)/q < t < (1 + 2b)/(1 - qb)$, an interval that is nonempty exactly when $(q + 2)a + qb > 1$.

[write] **6 October 2026; sources 34–39.** Source 35's Section I.61 extends the construction to polynomial phases of every degree, with exact thresholds $DC(\mathcal{S})$; at degree one its frontier is this one (its Corollary I.61.6).

Constants and ceilings at the boundary itself are governed by the finite inequalities, not by this strict-exponent statement. In particular the proposition establishes impossibility of any larger universal exponent; it does not prescribe an endpoint constant for a formula with exact powers and no rounding slack.

I.38.2 A small explicit obstruction

Take two phases, permitted lengths 4, 5, and tolerance $1/4$ for each. The construction gives

$$M_1 = M_2 = 11, \quad D_1 = 11, \quad D = 121, \quad C = 12, \quad R = 121 \cdot 11 = 1331.$$

The phases are $n/11$ and $n/121$. Flatness of a cell of length at least four requires each step phase to have distance at most $1/12$. If $11 \nmid d$, the first phase has distance at least $1/11$. If $11 \mid d$ but $121 \nmid d$,




Figure I.38.1: Two views of the partition theorem, with $q = 2$. Left: the prescribed-length and flexible-length scaling boundaries for target length R^a and diameter R^{-b} ; the dashed line is Gowers's relation $b = 2a$. Right: rigorous lower and upper bounds for $T_{\{L, L+1\}}(1/4, 1/4)/(16L^4)$ approach one. The shaded band bounds the unknown exact threshold; it is not numerical measurement of that threshold.

the second does. Therefore every cell stays in a residue class modulo 121. Every residue class has eleven points, while eleven is not $4u + 5v$ for nonnegative integers u, v .

For a prime-field witness, take permitted lengths 3, 4, tolerance $1/4$, $D = 49$, and $R = 245$. In $\mathbb{Z}/3433\mathbb{Z}$ use slopes 490 and 70. For $1 \leq d \leq 122$, the simultaneous necessary small-step inequalities allow only $d = 49, 98$. Both preserve the residues modulo 49, each of which contains five points. Five is not a sum of threes and fours. The accompanying script checks this claim with integer arithmetic, including primality and every candidate difference.

I.38.3 Why an exact prescribed number of cells is different

An exact number of nearly equal cells cannot be substituted silently for the allowed-length formulation. Here is a direct explanation of the repair already present in the repository.

Let $R = 4t + 2$ and request exactly $m = R/2 = 2t + 1$ nonempty cells with lengths differing by at most one. Their mean length is two, so every cell must have exactly two points. Choose a prime $p > 2R$ and the phase $\alpha = (p - 1)/(2p)$. Points of opposite parity in I_R have phase distance greater than $1/4$: for an odd difference s , the distance is $1/2 - |s|/(2p) > 1/4$. A flat cell of diameter at most $1/4$ is therefore parity-homogeneous. Each parity class has the odd cardinality $2t + 1$, so it cannot be partitioned into pairs. This remains a counterexample at arbitrarily large R .

I.39 Partial coverage and the exact cost of exceptions

Source 23, delivered Section 6.

The obstruction above is strong for exact partitions. Its behavior under deletion is different, and can also be computed exactly. This matters when a later argument only requires that most of the original points be retained.

For any finite allowed-length set $\mathcal{S}$, without a gcd assumption, define

$$A_{\mathcal{S}}(n) = \max\{s \leq n : s \in \langle \mathcal{S} \rangle\}, \quad g_{\mathcal{S}}(n) = n - A_{\mathcal{S}}(n).$$

The maximum exists because zero is representable.

Theorem I.39.1 (Exact packing for the rigid phase examples). *Choose radices M_i with $1/M_i > \varepsilon_i/(L-1)$ and phases $\alpha_i = 1/D_i$ as before. For every $R = Dt + r$, $0 \leq r < D$, the smallest number of uncovered points in a disjoint collection of flat progressions with lengths in $\mathcal{S}$ is exactly*

$$(D - r)g_{\mathcal{S}}(t) + rg_{\mathcal{S}}(t + 1). \quad (\text{I.39.1})$$

Proof. Every flat cell lies in a residue class modulo D . The residue chains have lengths t or $t + 1$, with multiplicities $D - r$ and r . A disjoint collection covers a representable number of points in each chain, so its uncovered count is at least the right side.

Conversely, cover the first $A_{\mathcal{S}}(n)$ points of each chain of length n by consecutive blocks whose lengths express that number in $\langle \mathcal{S} \rangle$. All blocks have step D , on which all phases are exactly constant. This attains the lower bound. $\square$

Corollary I.39.2 (Exact threshold of the constructed phase vector). *For the mixed-radix phases in Theorem I.39.1, if $\gcd(\mathcal{S}) = 1$, the eventual exact-partition threshold is exactly $DC(\mathcal{S})$.*

Proof. At $R = DC - 1$ the last residue chain has the gap cardinality $C - 1$. At every $R \geq DC$, all residue-chain cardinalities are at least C , and hence representable. Theorem I.39.1 then gives zero uncovered points. $\square$

For $\mathcal{S} = \{L, L + 1\}$, the last gap $F = L^2 - L - 1$ has $g_{\mathcal{S}}(F) = 1$, because $F - 1 = (L - 2)(L + 1)$. Thus the exact-partition obstruction at $R = DF$ disappears after deleting exactly D points, a fraction $1/F$. The conductor obstruction alone therefore cannot exclude an almost partition with exceptional proportion appreciably larger than L^{-2} .

Proposition I.39.3 (A general almost-partition into one exact length). *Given $L \geq 2$, arbitrary affine phases and positive tolerances, put*

$$Q = \prod_{i=1}^q \left\lceil \frac{L - 1}{\varepsilon_i} \right\rceil.$$

There is a disjoint collection of flat progressions of length exactly L covering all but at most $(L - 1)Q$ points of I_R . In particular, an exceptional proportion at most $\eta > 0$ is guaranteed when

$$R \geq \frac{(L - 1)Q}{\eta}.$$

Proof. Choose $d \leq Q$ by Lemma I.36.3. Cut each residue chain into blocks of length L , discarding a remainder of at most $L - 1$. There are at most d nonempty chains, and each block is flat by the same phase-span estimate as before. $\square$

Corollary I.39.4 (Sharp scale when a small exceptional set is allowed). *Let $T_{L,\eta}(\varepsilon)$ be the least eventual threshold for covering at least $(1 - \eta)R$ points by disjoint flat progressions of length exactly L . For fixed q , as $L \rightarrow \infty$ and $\eta \rightarrow 0^+$,*

$$T_{L,\eta}(\varepsilon) \sim \frac{L^{q+1}}{\eta \varepsilon_1 \cdots \varepsilon_q}, \quad (\text{I.39.2})$$

uniformly for $0 < \varepsilon_i < 1/2$.

Proof. The upper bound follows from Proposition I.39.3 and the uniform ceiling estimates used earlier. For the lower bound take the mixed-radix phases with $M_i = \lceil (L-1)/\varepsilon_i \rceil - 1$ and $D = \prod_i M_i$. Choose $t = kL - 1$ to be the largest integer of that form strictly smaller than $(L-1)/\eta$. For sufficiently small η , it is positive and

$$t = \frac{L-1}{\eta} + O(L), \quad \frac{L-1}{t} > \eta.$$

At $R = Dt$, Theorem I.39.1 with $\mathcal{S} = \{L\}$ gives uncovered fraction $(L-1)/t$, strictly greater than η . Consequently $T_{L,\eta} > Dt$. Now

$$Dt = (1 + O_q(L^{-1}) + O(\eta)) \frac{L^{q+1}}{\eta \prod_i \varepsilon_i},$$

which squeezes the threshold to the asserted asymptotic. $\square$

The exact-cover scale and the almost-cover scale coincide at the level of powers when η is of order L^{-1} . This comparison uses a universal bound over all phases. The particular last-gap example is easier to repair, as the exact formula (I.39.1) demonstrates.

I.39.1 Constructive use and computational complexity

The upper theorem is constructive at the finite combinatorial level. For rational phases and tolerances, enumerate $d = 1, \dots, B$, where $B = \lceil (H-1)^q / (\prod_i \varepsilon_i) \rceil - 1$, testing $\|d\alpha_i\|_{\mathbb{T}} \leq \varepsilon_i / (H-1)$ for every i . All comparisons can be made exactly. Lemma I.36.2 guarantees a successful difference in this range, so the stated sharp threshold leaves every residue chain at least C points long. Then compute one representation in $\langle \mathcal{S} \rangle$ for each of the two possible residue-chain lengths. Dynamic programming up to the largest chain length suffices; for an interval of allowed lengths, the explicit distribution of excess in Proposition I.37.3 gives a simpler construction. Cutting and listing all cells takes time linear in the output size.

This elementary implementation is not presented as an optimal algorithm for simultaneous Diophantine approximation. It exposes the exact certificates needed by a checker: one integer difference, its phase increments, the chain cardinalities, and representations of those cardinalities as sums of permitted lengths.

I.40 Corollary 5.8 without a scale threshold: degrees one and two [write]

Added in the fifth write, 6 October 2026.

This chapter is about Gowers’s Corollary 5.8, the local density increment that turns correlations with polynomial phases on a cover by progressions into a progression of higher density. Source 02’s Section I.10 gives a replacement interface and declines to restate the corollary itself (“assertions about the length ... must be supplied by an actual partition theorem”); the formal statement `corollary_5_8` (Section05.lean, line 242, at HEAD 2dfe46f44) is open, and of the corollary itself the repository proves only the scale-qualified repair `corollary_5_8_with_scale_holds` (Proofs05FullPartition.lean, line 94). It also proves, since 1c1371304, the refinement step with cell-dependent phases, `variable_polynomial_phase_refinement` (Proofs05VariablePhas

`eRefinement.lean`, line 13: common degree, coefficients varying between cells, half the correlation kept); that is an ingredient, not the density increment. The next four sections are source 26's Sections 2–4 and Appendix A, which prove the degree-one case with no scale hypothesis on every cyclic group. The chapter is planned for more sources: source 32 (Report276, placed by `c5513046e`, write pending), the degree-two sibling of source 26 with the same pin and format, belongs directly after it, and its placement records it as proving the degree-two case; sources 41 (Report277) and 42 (Report278), placed by `2dfe46f44` while this write was in progress, follow it, and their placement records 41 as a written proof of the whole statement in every degree with no scale premise, extending source 32, and 42 as replacing source 26's cube-root length cap by a square root. Those are for the writes that add them. Source 26's introduction is Section F.31; its source comparison, questions and companion are in Part V (Sections V.105–V.106). Its notation is in the table at the head of this Part.

What is proved, in one paragraph. Let $A \subseteq \mathbb{Z}/N\mathbb{Z}$ be covered by disjoint proper modular progressions of shortest length r , with affine phases whose summed correlations with $\mathbf{1}_A - \delta$ are at least αN . Then some proper progression has density at least $\delta + \alpha/3$ and length at least $\min\{(r(1 - \delta)/(128\mu))^{1/3}, (N/8)^{1/3}\}$, where μ is the normalized balanced mass of the cover (Theorem I.41.1), and some progression inside a parent has density $\delta + c\alpha$ and length $\sqrt{(1 - 2c)\alpha r/(8\pi\mu)}$ (Theorem I.41.2); for two-comparable parents this gives $\max\{(N/M)^{1/3}/8, \sqrt{\alpha(N/M)/(96\pi(1 - \delta))}\}$ (Corollary I.41.3). The two regimes are a second-moment increment over all steps, with the nonproper steps removed, when α is small (Lemma I.42.1), and phase-flat refinement with positive-cell extraction when it is large (Lemma I.42.7). Appendix A converts proper modular boxes into an almost-cover by ordinary integer boxes with one common step (Theorem I.44.1), and shows that exact long ordinary partitions of arbitrary proper modular parents are impossible. The source credits the cube-root scale and the mechanisms to Gowers, Zhao and Lewko, and claims no higher-degree result.

Relation to the formal statements (checked at HEAD `2dfe46f44`). With $k = 1$, `PolynomialOn` (`Definitions.lean`, line 222) is exactly an affine map and `polynomialPartitionConstant 1 = 16` (`Section05.lean`, line 45), so the degree-one instance of `corollary_5_8` asks for gain $\alpha/16$ at length $(N/M)^{1/16}/8$; Corollary I.41.3 implies it on paper, and that instance holds trivially when $N/M \leq 2^{48}$ (note in Section F.31). The ledger is unchanged: the formal statement quantifies over every degree. A formalization candidate, not formalized. Appendix A bears on the caveat of `FORMALIZATION_STATUS.txt` that proper boxes may wrap in integer representatives; it gives an almost-cover, not the exact-partition and base-loss interface of Section 16 (the source says so), and sits beside source 23's Section II.103 and source 19's Corollary II.70.8.

[write] **6 October 2026; sources 28–33.** *Source 32 is now printed here*, as planned above: its Sections 2–5 follow source 26's Appendix A at the end of this chapter (Sections I.45–I.48; introduction Section F.37, source comparison and questions Sections V.125–V.126). At degree two, with no scale hypothesis and over every cyclic group, some proper progression has density at least $\delta + \alpha/3$ and length at least $\max\{1, (N/M)^{1/39}/17\}$ (Theorem F.37.1), through an explicit recurrence $p \leq 2^{10}R^5$ (Lemma I.45.3), lossless quadratic partitions above $2^{116}L^{38}$ (Lemma I.46.2) and source 26's second-moment step. It implies the degree-two instance of `corollary_5_8`, so the instances $k \leq 2$ have written proofs; the statement stays open in the ledger (checked at HEAD `5d4718772`). Sources 41 (all degrees, extending source 32) and 42 (square-root caps) are placed and not yet written.

[write] **7 October 2026; sources 40–46.** *Sources 41 and 42 are now printed*, at the end of Part I (Section I.62), because this chapter is no longer the last of Part I. Source 41 proves the corollary in every degree with no scale premise (gain $\alpha/3$, length $\max\{1, (N/M)^{1/D_k}/8\}$), keeping this chapter's bounds in degrees one and two; its Sections 2 and 5 are source 32's Sections 2 and 4 verbatim. Source 42 replaces the scale $N \geq L^3$ of the variance step by $N \geq 2(L - 1)^2 + 1$ and, in source 26's affine bound, the cap $(N/8)^{1/3}$ by $\sqrt{N/2}$. *The formal corollary_5_8 is now exact* at HEAD `6d4942222`: `corollary_5_8_holds`

(Proofs05ThresholdFreeLocalization.lean, line 86), committed in b548d998e while the eighth write was in progress, with this chapter’s two-regime architecture and the repository’s own finite localization; the degree-one and degree-two cases are `corollary_5_8_degree_one` and `corollary_5_8_degree_two` (Section I.62).

I.41 Definitions and the two main bounds

Source 26, delivered Section 2.

Let $N \geq 2$ be an integer and write $G = \mathbb{Z}/N\mathbb{Z}$. For a residue y define

$$e_N(y) = \exp(2\pi i \tilde{y}/N),$$

where $\tilde{y}$ is any integer representative. A *proper modular progression* of length $n \geq 1$ is a set

$$P = \{u + jd : 0 \leq j < n\} \subseteq G$$

with this parameterization injective. The step d need not be a unit. Singletons are allowed. In particular, “proper” does not mean that standard integer representatives form a nonwrapping progression.

Let $A \subseteq G$, let $\delta = |A|/N$, and let $f = \mathbf{1}_A - \delta$. Suppose $P_1, \dots, P_M$ are nonempty pairwise disjoint proper progressions, $M \geq 1$, and

$$A \subseteq U := \bigcup_{i=1}^M P_i.$$

For each i let $\phi_i(x) = a_i x + b_i$ be an affine map $G \rightarrow G$. Our correlation hypothesis is

$$\sum_{i=1}^M \left| \sum_{x \in P_i} f(x) e_N(-\phi_i(x)) \right| \geq \alpha N, \quad \alpha > 0. \quad (\text{I.41.1})$$

Define the cover parameters

$$\tau = \frac{|U|}{N}, \quad r = \min_i |P_i|, \quad \mu = \frac{1}{N} \sum_{x \in U} |f(x)|. \quad (\text{I.41.2})$$

The hypotheses force $0 < \delta < 1$, and direct counting gives

$$\mu = \delta(1 + \tau - 2\delta), \quad (\text{I.41.3})$$

$$0 < \alpha \leq \mu \leq 2\delta(1 - \delta), \quad \mu \geq \delta(1 - \delta), \quad \delta \leq \tau \leq 1, \quad (\text{I.41.4})$$

$$\sum_{x \in U} f(x) = \delta N(1 - \tau) \geq 0. \quad (\text{I.41.5})$$

For example, (I.41.3) follows by separating the δN points of A , where $|f| = 1 - \delta$, from the $(\tau - \delta)N$ other points of U , where $|f| = \delta$. The first inequality in (I.41.4) follows from (I.41.1) and the triangle inequality. If δ were zero or one, every correlation would vanish.

Theorem I.41.1 (Threshold free cube root bound). *Under the hypotheses above, there is a proper modular progression Q such that*

$$|Q| \geq \min \left\{ \left(\frac{r(1 - \delta)}{128\mu} \right)^{1/3}, (N/8)^{1/3} \right\}, \quad \text{dens}_Q(A) := \frac{|A \cap Q|}{|Q|} \geq \delta + \frac{\alpha}{3}. \quad (\text{I.41.6})$$

No scale threshold or lower bound on α is imposed, and no primality assumption is needed. The progression Q need not lie in U .

[write] **7 October 2026; sources 40–46.** Source 42’s Theorem I.74.1 replaces the modulus cap $(N/8)^{1/3}$ by $\sqrt{N/2}$, never smaller; the term $(r(1-\delta)/(128\mu))^{1/3}$ is unchanged (note there).

Theorem I.41.2 (Correlation dependent square root bound). *Under the same hypotheses, for every real c with $0 < c < 1/2$ there is a proper modular progression Q , contained in one of the parents P_i , such that*

$$|Q| \geq \sqrt{\frac{(1-2c)\alpha r}{8\pi\mu}}, \quad \text{dens}_Q(A) \geq \delta + c\alpha. \quad (\text{I.41.7})$$

This result has no separate modulus-dependent cap. In particular, $c = 1/3$ gives length at least $\sqrt{\alpha r/(24\pi\mu)}$ and gain $\alpha/3$.

The two conclusions can be combined at gain $\alpha/3$: select the progression supplied by the theorem with the larger length guarantee. This gives the maximum of the two right-hand sides, but parent containment is assured only when the square-root conclusion is selected. The square-root estimate depends on α ; it is not an α -independent replacement of the cube-root estimate.

Corollary I.41.3 (Intrinsic cover and comparable parents). *The cube-root guarantee in Theorem I.41.1 is at least*

$$\left(\frac{r}{256\tau}\right)^{1/3}. \quad (\text{I.41.8})$$

If, in addition, $|P_i| \leq B|P_j|$ for all i, j , where $B \geq 1$, put $t = N/M$. There are progressions with gain $\alpha/3$ and respective length guarantees

$$\left(\frac{t}{256B}\right)^{1/3}, \quad \sqrt{\frac{\alpha t}{48\pi B(1-\delta)}}. \quad (\text{I.41.9})$$

For $B = 2$ and every $N \geq 2$, their maximum is

$$\max \left\{ \frac{(N/M)^{1/3}}{8}, \sqrt{\frac{\alpha(N/M)}{96\pi(1-\delta)}} \right\}. \quad (\text{I.41.10})$$

Proof. By (I.41.4), $\mu \leq 2\tau(1-\delta)$, while $r \leq |U| = \tau N$. Thus $r/(256\tau) \leq N/256 \leq N/8$, proving (I.41.8). The same mass bound gives the square-root guarantee $\sqrt{\alpha r/(48\pi\tau(1-\delta))}$. Comparability gives $|U| \leq BMr$, hence $r/\tau \geq N/(BM) = t/B$. $\square$

[write] **6 October 2026; sources 21–27.** *It implies the degree-one instance of `corollary_5_8` (checked at the write).* The formal hypotheses at $k = 1$ are those of Section I.41 with $B = 2$: `PolynomialOn 1` (`Definitions.lean`, line 222) is an affine map $c_0 + c_1x$, `ModAP.IsProper` is injective parameterization, the parents are disjoint and two-comparable, and the correlation uses `balanced` and `exponential`, that is $\mathbf{1}_A - \delta$ and e_N . The parents are nonempty (one empty parent would make all of them empty, and then $A = \emptyset$, contradicting $\alpha > 0$), so $M \leq N$ and $(N/M)^{1/3}/8 \geq (N/M)^{1/16}/8$; and $\delta + \alpha/3 \geq \delta + \alpha/16$. So (I.41.10) gives the formal conclusion with room to spare, for every modulus $N \geq 2$ (and $N = 1$ is vacuous, Section V.105). Not formalized.

[write] **6 October 2026; sources 28–33.** The degree-two analogue is source 32’s Theorem F.37.1: gain $\alpha/3$ at length $\max\{1, (N/M)^{1/39}/17\}$, which implies the degree-two instance of `corollary_5_8`; together the two give written proofs of the instances $k \leq 2$ (note there).

I.42 Two elementary mechanisms

Source 26, delivered Section 3.

I.42.1 Moments and the removal of nonproper directions

Lemma I.42.1 (Proper progression increment over every cyclic modulus). *Let $0 < \delta < 1$, and let $L \geq 2$ be an integer with $N \geq L^3$. Some proper L -term modular progression has A -density at least*

$$\delta + \frac{1 - \delta}{2L}. \quad (\text{I.42.1})$$

Proof. Choose a, d independently and uniformly from all of G , including $d = 0$, and put

$$X(a, d) = \frac{1}{L} \sum_{j=0}^{L-1} \mathbf{1}_A(a + jd).$$

Here repeated positions are counted with their index multiplicity. For every fixed d , averaging over a gives $\mathbb{E}_a X = \delta$.

For distinct indices j, k , put $H = (k - j)G$, a subgroup of G . Under the bijective change of variables $x = a + jd$, the pair (x, d) is still uniform on G^2 , and $(k - j)d$ is uniform on H . Consequently

$$\mathbb{E}[\mathbf{1}_A(a + jd)\mathbf{1}_A(a + kd)] = \sum_{C \in G/H} \frac{|H|}{N} \left(\frac{|A \cap C|}{|H|} \right)^2 \geq \delta^2. \quad (\text{I.42.2})$$

Indeed the densities in the cosets have weighted mean δ , and their weighted squared deviations from that mean are nonnegative. Thus every off-diagonal covariance is nonnegative. The L diagonal terms give

$$\text{Var}_{a,d} X \geq \frac{\delta(1 - \delta)}{L}. \quad (\text{I.42.3})$$

Let D_{bad} be the directions for which the displayed progression is not proper. Properness depends only on d . Such a direction satisfies $hd = 0$ for some integer $1 \leq h < L$, and conversely this condition produces a repetition. The kernel of multiplication by h on G has $\gcd(h, N) \leq h$ elements. To see the equality, write $h = gh'$, $N = gN'$ with $g = \gcd(h, N)$ and $\gcd(h', N') = 1$; the condition $N \mid hd$ is equivalent to $N' \mid d$, leaving exactly g residues. The union bound now yields

$$b := \frac{|D_{\text{bad}}|}{N} \leq \frac{1}{N} \sum_{h=1}^{L-1} h = \frac{L(L-1)}{2N} < \frac{1}{2L} < 1. \quad (\text{I.42.4})$$

Let $h_* = \max(X(a, d) - \delta)$ over proper directions. Such directions exist, since $d = 1$ is proper when $L \leq N$. For a proper direction, $0 \leq X \leq \delta + h_*$ and $\mathbb{E}_a X = \delta$, so

$$\text{Var}_a X \leq \delta h_*.$$

For a bad direction, $0 \leq X \leq 1$ instead gives $\text{Var}_a X \leq \delta(1 - \delta)$. Since the conditional means are all δ , the total variance is exactly the average of the conditional variances. Combining with (I.42.3) gives

$$\frac{\delta(1 - \delta)}{L} \leq \text{Var}_{a,d} X \leq (1 - b)\delta h_* + b\delta(1 - \delta).$$

It follows that

$$h_* \geq (1 - \delta) \frac{1/L - b}{1 - b} \geq \frac{1 - \delta}{2L}. \quad (\text{I.42.5})$$

For the last inequality, (I.42.4) gives $1/L - b \geq 1/(2L)$ and $1 - b \leq 1$. On proper directions X is the actual set density, proving the lemma. $\square$

[write] **7 October 2026; sources 40–46.** *Two later facts* (sources 41 and 42; checked at HEAD 6d4942222). The lemma is repeated verbatim as source 32’s Lemma I.47.1 and source 41’s Lemma I.66.1, and is formalized as `exists_proper_progression_density_increment` (`Proofs05VarianceIncrement.lean`, line 89). Source 42 proves the same gain under $N \geq 2(L - 1)^2 + 1$ (Corollary F.47.2), with an arithmetic refinement and a sharp constant.

Remark I.42.2 (The exact prime and short index identity). A useful exact diagnostic uses a uniform on G and d uniform on its nonzero residues. Write $\text{spf}(N)$ for the least prime divisor of N . If $1 \leq L \leq \text{spf}(N)$, every nonzero index difference $j - k$ is a unit. Every progression is then proper, and for $j \neq k$ the map

$$(a, d) \mapsto (a + jd, a + kd)$$

is a bijection onto the ordered distinct residue pairs: its inverse is $d = (j - k)^{-1}(x - y)$, $a = x - jd$. Expanding X^2 therefore gives

$$\begin{aligned} \mathbb{E}X^2 &= \frac{\delta}{L} + \frac{L - 1}{L} \frac{\delta(\delta N - 1)}{N - 1}, \\ \text{Var } X &= \frac{\delta(1 - \delta)(N - L)}{L(N - 1)}. \end{aligned}$$

This equality is not valid without the index restriction. For example, $N = 4$, $L = 3$, $A = \{0, 2\}$ gives actual variance $11/108$, whereas the displayed formula would give $1/36$. The all-direction lower bound (I.42.3), followed by removal of bad directions, is what removes any prime-factor restriction from the main result. Sampling only unit steps would not justify the same pair-uniformity argument over a composite modulus.

I.42.2 Phase flat residue cells with exact rounding

Lemma I.42.3 (Elementary circle gap). *For a real number θ and an integer $R \geq 1$, there exist integers q, v with*

$$1 \leq q \leq R, \quad |q\theta - v| \leq \frac{1}{R + 1}.$$

Proof. If two of $0, \theta, \dots, R\theta$ agree modulo one, their index difference gives the assertion with zero error. Otherwise these $R + 1$ points are distinct on the circle. Their cyclic gaps sum to one, so a gap has length at most $1/(R + 1)$. The index difference of its endpoints has absolute value between one and R ; taking its positive absolute value changes at most the sign of the error. The wraparound gap is treated the same way modulo an integer. $\square$

Lemma I.42.4 (Balanced chunks). *If integers n, L satisfy $n \geq L \geq 1$, a chain of n points can be partitioned consecutively into chunks with cardinalities in $[L, 2L - 1]$.*

Proof. Write $n = hL + b$ with integers $h \geq 1$ and $0 \leq b < L$. Use $h - 1$ chunks of size L and one chunk of size $L + b$. There is no discarded tail, including when $b = 0$. $\square$

Lemma I.42.5 (Phase flat refinement). *Let $P = \{u + jd : 0 \leq j < n\}$ be a proper progression, let $\phi(x) = ax + b$ be affine, and let $1 \leq L \leq n$ be an integer. There is a disjoint partition of P into proper progressions C with*

$$L \leq |C| \leq 2L - 1,$$

and unit complex numbers z_C such that

$$|e_N(-\phi(x)) - z_C| \leq \frac{2\pi L^2}{n} \quad (x \in C). \quad (\text{I.42.6})$$

This holds over every cyclic modulus.

Proof. Choose a real representative θ of ad/N and put $R = \lfloor n/L \rfloor \geq 1$. Lemma I.42.3 supplies $1 \leq q \leq R$ and

$$q\theta = v + \varepsilon, \quad v \in \mathbb{Z}, \quad |\varepsilon| \leq \frac{1}{R+1}.$$

Partition $\{0, \dots, n-1\}$ into residue classes modulo q , in increasing order in each class. Each chain has at least $\lfloor n/q \rfloor \geq L$ points. Apply Lemma I.42.4 to every chain and take the images in P . They are disjoint and proper because the entire parent parameterization is injective.

For one resulting cell write $C = \{x_0 + jqd : 0 \leq j < m\}$, where $L \leq m \leq 2L - 1$, and define

$$z_C = e_N(-\phi(x_0)) \exp\left(-2\pi i \varepsilon \frac{m-1}{2}\right). \quad (\text{I.42.7})$$

For integer j , the integer part vj disappears, so

$$e_N(-\phi(x_0 + jqd)) = e_N(-\phi(x_0)) \exp(-2\pi i \varepsilon j).$$

The chord-arc inequality $|e^{iu} - e^{iv}| \leq |u - v|$ now gives

$$\begin{aligned} |e_N(-\phi(x_0 + jqd)) - z_C| &\leq 2\pi |\varepsilon| \left| j - \frac{m-1}{2} \right| \\ &\leq \frac{2\pi(L-1)}{R+1} < \frac{2\pi L^2}{n}. \end{aligned}$$

Here $R+1 > n/L$; for $L=1$ the left side is zero. The weak bound (I.42.6) follows. $\square$

Remark I.42.6 (The midpoint must use the reduced increment). The midpoint in (I.42.7) is formed after writing $q\theta = v + \varepsilon$. Replacing ε by the unreduced $q\theta$ multiplies it by $\exp(-\pi i v(m-1))$, which can be -1 . For an exact regression, take $N = 6$, $\phi(x) = 3x$, and the proper two-point parent $P = \{0, 2\}$ of step 2, with target $L = 2$. Using the integer lifts gives $\theta = 3 \cdot 2/6 = 1$, $R = 1$, $q = 1$, $v = 1$ and $\varepsilon = 0$. Both endpoint phases and the reduced midpoint equal $+1$. The unreduced midpoint is $\exp(-\pi i) = -1$, at distance 2 from both endpoints, although the reduced-lift radius is zero. Thus unreduced midpoint wording would be false even for a phase constant on the entire parent.

I.42.3 Turning correlation into positive balanced mass

Lemma I.42.7 (Refinement transfer). *Under the hypotheses of Section I.41, let $1 \leq L \leq r$ be an integer. There is a partition of U into parent-contained proper progression cells, each of cardinality in $[L, 2L - 1]$, such that, with $F_C = \sum_{x \in C} f(x)$,*

$$\sum_C |F_C| \geq \alpha N - E, \quad E = \frac{2\pi L^2 \mu N}{r}. \quad (\text{I.42.8})$$

If $E \leq (1 - 2c)\alpha N$ for $0 < c < 1/2$, one of these cells has density at least $\delta + c\alpha$.

Proof. Apply Lemma I.42.5 separately to each parent. For that parent, subtract the constant phase z_C on each child and use the triangle inequality. Summing gives

$$\sum_i \left| \sum_{P_i} f(x) e_N(-\phi_i(x)) \right| \leq \sum_C |F_C| + \frac{2\pi L^2}{r} \sum_U |f(x)|.$$

This proves (I.42.8). The F_C are real and (I.41.5) gives $\sum_C F_C \geq 0$. Therefore

$$\sum_{F_C > 0} F_C = \frac{1}{2} \left(\sum_C |F_C| + \sum_C F_C \right) \geq \frac{\alpha N - E}{2}. \quad (\text{I.42.9})$$

If $E \leq (1 - 2c)\alpha N$, the right side is at least $c\alpha N > 0$. The positive cells together have at most N points, so at least one satisfies $F_C/|C| \geq c\alpha$. $\square$

Only $A \subseteq U$ was assumed. In general $\sum_U f$ is nonnegative rather than zero; the identity (I.41.5) is exactly what is needed in (I.42.9).

I.43 Proofs of the main theorems

Source 26, delivered Section 4.

I.43.1 The cube root bound

Proof of Theorem I.41.1. Put

$$H = \frac{r(1 - \delta)}{\mu}, \quad x = \min\{(H/128)^{1/3}, (N/8)^{1/3}\}, \quad L = \lceil x \rceil.$$

The integer $L \geq 1$ satisfies $L \geq x$. If $L = 1$, take a singleton in A . Its density is one, which exceeds $\delta + \alpha/3$ since $\alpha \leq 2\delta(1 - \delta)$ and $0 < \delta < 1$.

Assume $L \geq 2$. Then $x > 1$, and therefore

$$L = \lceil x \rceil < 2x, \quad H > 16L^3, \quad N > L^3. \quad (\text{I.43.1})$$

The last two inequalities follow from the respective upper bounds defining x . Lemma I.42.1 is therefore applicable.

If

$$\alpha \leq \frac{3(1 - \delta)}{2L}, \quad (\text{I.43.2})$$

that lemma supplies a proper progression of length exactly L and density at least $\delta + \alpha/3$. This settles the small-correlation range without refining any parent.

Otherwise $\alpha > 3(1 - \delta)/(2L)$. Since $\alpha \leq \mu$, (I.43.1) gives

$$r = \frac{H\mu}{1 - \delta} > 16L^3 \frac{3}{2L} = 24L^2 > L.$$

Refinement is therefore legal. Its error in Lemma I.42.7 satisfies

$$\begin{aligned} E &= \frac{2\pi L^2 \mu N}{r} = \frac{2\pi L^2 (1 - \delta) N}{H} \\ &< \frac{\pi(1 - \delta) N}{8L} < \frac{(1 - \delta) N}{2L} < \frac{\alpha N}{3}, \end{aligned}$$

using $\pi < 4$. Lemma I.42.7, or its strict version from (I.42.9), yields a parent-contained cell of length at least L with density exceeding $\delta + \alpha/3$. This completes both cases. $\square$

Remark I.43.1 (A complementary short index refinement). The exact identity in Remark I.42.2 also retains the bound

$$|Q| \geq \min \left\{ \text{spf}(N), \left(\frac{r(1 - \delta)}{128\mu} \right)^{1/3} \right\}, \quad \text{dens}_Q(A) \geq \delta + \alpha/3.$$

Here is the rounding and range check. Put $H = r(1 - \delta)/\mu$, $x_0 = (H/128)^{1/3}$ and $L = \min\{\text{spf}(N), \lceil x_0 \rceil\}$. The singleton case is as before. Otherwise $L < 2x_0$, hence $H > 16L^3$. Also $H \leq r/\delta \leq N^2$, since $\mu \geq \delta(1 - \delta)$ and the nonempty A has $\delta \geq 1/N$. Therefore $N > 2L$. For nonzero-step sampling the exact identity and $\text{Var } X \leq \delta \max(X - \delta)$ give gain at least $(1 - \delta)/(2L)$. The two cases of the preceding proof now apply unchanged. For prime N the least-prime-factor cap never binds, so this gives the uncapped exact-mass cube root. This complementary bound may be stronger than Theorem I.41.1 for sparse sets with prime or rough modulus; the all-modulus theorem is not claimed to dominate it pointwise.

The exact-mass form exposes the mechanism particularly directly: the large- α condition itself forces $r > 24L^2$. For the comparable-parent formulation one can instead track

$$r \geq \frac{\tau t}{B}, \quad \mu \leq 2\delta(1 - \delta), \quad \delta \leq \tau.$$

Then the freezing error is at most

$$\frac{4\pi B L^2 \delta(1 - \delta) N}{\tau t} \leq \frac{4\pi B L^2 (1 - \delta) N}{t}. \quad (\text{I.43.3})$$

This is the δ/τ cancellation. It explains why sparse covers do not introduce an unaccounted inverse-density scale assumption.

I.43.2 The square root bound

Proof of Theorem I.41.2. Set

$$R = \sqrt{\frac{(1 - 2c)\alpha r}{2\pi\mu}}, \quad L = \max\{1, \lfloor R \rfloor\}.$$

If $L = 1$, then $R < 2$ and a singleton in A has length greater than $R/2$, the required right side of (I.41.7). It lies in a parent because $A \subseteq U$. Its density is at least $\delta + c\alpha$: indeed

$$c\alpha < \alpha/2 \leq \delta(1 - \delta) < 1 - \delta.$$

Suppose $L \geq 2$. Since $\alpha \leq \mu$ and $0 < 1 - 2c < 1$,

$$R^2 \leq \frac{(1 - 2c)r}{2\pi} < r.$$

Thus $r > R^2 \geq L^2 \geq L$, and Lemma I.42.7 applies. Moreover

$$E = \frac{2\pi L^2 \mu N}{r} \leq \frac{2\pi R^2 \mu N}{r} = (1 - 2c)\alpha N.$$

The selected cell has density at least $\delta + c\alpha$. Finally, $R \geq 2$ implies $L = \lfloor R \rfloor \geq R/2$, proving the length bound with the stated constant. No moment estimate or primality assumption was used. $\square$

For two-comparable parents, the simplified square-root bound in (I.41.10) is at least as large as the simplified cube-root bound precisely when

$$\frac{\alpha}{1 - \delta} \geq \frac{3\pi}{2} t^{-1/3}.$$

This comparison concerns the two convenient lower bounds, not the optimal length obtainable from the hypotheses. The exact-mass bounds can be stronger.

I.44 Common step integer box rectification

Source 26, delivered Appendix A.

The width of a nonempty box is its smallest axis cardinality. A common modular step across the axes of one proper box can be converted into a common positive integer step on a large retained product. Loss is measured in the ambient product. No density-increment property of an arbitrary retained subset is inferred merely from its size.

Theorem I.44.1 (Ordinary box almost cover). *Let $N \geq 1$, $D \geq 1$, and*

$$P = P_1 \times \cdots \times P_D \subseteq (\mathbb{Z}/N\mathbb{Z})^D, \quad P_i = \{a_i + jd : 0 \leq j < n_i\},$$

where every axis parameterization is injective and $m = \min_i n_i \geq 1$. For $0 < \eta \leq 1$, set the integer

$$L = 1 + \left\lfloor \frac{\eta\sqrt{m}}{3D} \right\rfloor.$$

There is $G \subseteq P$ with $|P \setminus G| \leq \eta|P|$ and a disjoint partition of G into boxes whose axes, in standard representatives $0, \dots, N - 1$, are ordinary arithmetic progressions. Every axis has cardinality in $[L, 2L - 1]$, and all axes of all output boxes use one positive integer difference Δ . In particular the width is at least $L > \eta\sqrt{m}/(3D)$.

Proof. If $L = 1$, use singleton boxes and $\Delta = 1$. Suppose $L \geq 2$. Then $m \geq 2$, and $Q = \lfloor \sqrt{m} \rfloor$ satisfies $1 \leq Q < m$. Lemma I.42.3, applied to d/N , gives $1 \leq q \leq Q$ whose centered residue b for qd has $|b| \leq N/(Q + 1)$. It is nonzero: otherwise indices 0 and $q < m \leq n_i$ would have the same parent image. Put $\Delta = |b| > 0$.

Partition every parent index interval into its q residue chains. If $b < 0$, reverse every chain, so that all traversals have modular increment $+\Delta$. Consecutive standard representatives then differ by

either Δ or $\Delta - N$. Cut at every decrease of standard representatives. Each resulting run is an ordinary progression of positive difference Δ .

For a chain of h points with initial representative $z \in [0, N - 1]$, the number of cuts is exactly

$$W = \left\lfloor \frac{z + (h - 1)\Delta}{N} \right\rfloor < \frac{(h - 1)\Delta}{N} + 1.$$

Summing the number $W + 1$ of runs over all q chains on axis i gives a run count R_i satisfying

$$R_i \leq \frac{n_i}{Q + 1} + 2q \leq \frac{n_i}{Q + 1} + 2Q.$$

Delete the runs of cardinality less than L . Their total cardinality D_i is at most $(L - 1)R_i$, and therefore

$$\begin{aligned} \frac{D_i}{n_i} &\leq (L - 1) \left(\frac{1}{Q + 1} + \frac{2Q}{n_i} \right) \\ &\leq \frac{3(L - 1)}{\sqrt{m}} \leq \frac{\eta}{D}. \end{aligned}$$

Here $Q + 1 > \sqrt{m}$, $Q \leq \sqrt{m}$, and $n_i \geq m$. Split each surviving run by Lemma I.42.4, without further loss. Let G_i be the union of the surviving chunks and $G = \prod_i G_i$. Products of chunks partition G disjointly, retain the same Δ , and satisfy

$$\frac{|P \setminus G|}{|P|} = 1 - \prod_i \left(1 - \frac{D_i}{n_i} \right) \leq \sum_i \frac{D_i}{n_i} \leq \eta.$$

The product inequality follows by repeated use of $(1 - u)(1 - v) \geq 1 - u - v$ for $u, v \in [0, 1]$. $\square$

[write] **6 October 2026; sources 21–27.** *Relation to the Section 16 material.* The boxes here are proper modular boxes made into ordinary integer boxes on most of their mass; source 23's Section II.103 works with proper modular boxes throughout and lists integer-representative fidelity among its non-claims, and source 19's Corollary II.70.8 needs an inherited-coordinate hypothesis. As the source says in Section I.44.1, this almost-cover does not give the exact-partition and base-loss interface of the formal Section 16 statements, so it bears on, and does not discharge, the caveat of `FORMALIZATION_STATUS.txt` that proper boxes may wrap in integer representatives.

The proof handles composite moduli and nonunit original steps. If the common step is zero, proper nonempty axes have length one, so the singleton branch applies. It also handles $N = 1$. Although the weak estimate at $\eta = 1$ permits zero retention, the construction has nonempty retention: in the nonsingleton branch the bound $1/(Q + 1) + 2Q/n_i < 3/\sqrt{m}$ is strict.

Corollary I.44.2 (Anisotropic budget). *For integers $1 \leq Q < m$ and $L \geq 1$, the same conclusion with axis lengths in $[L, 2L - 1]$ holds whenever*

$$(L - 1) \sum_{i=1}^D \left(\frac{1}{Q + 1} + \frac{2Q}{n_i} \right) \leq \eta.$$

Proof. Repeat the proof, sum its individual deletion bounds, and use the product inequality. For $L = 1$ the singleton partition suffices. $\square$

The integer hypotheses are essential: cardinality intervals need not contain an integer for an arbitrary real L . For numerical optimization of the budget, write $S = \sum_i 1/n_i$. The real function $D/(Q + 1) + 2SQ$ has its stationary point at $\sqrt{D/(2S)} - 1$. Clamping this to $[1, m - 1]$ and checking neighboring integers and endpoints gives the discrete optimum by convexity.

I.44.1 What this gives for the source geometry

Apply Theorem I.44.1 separately to a disjoint family of proper modular boxes of width at least w . The union of the retained products loses at most η of the ambient product mass; each output box has width greater than $\eta\sqrt{w}/(3D)$. The positive integer step is common within each original box's descendants. It can differ across original boxes. Existing identities expressing agreement with a fixed affine or multilinear function remain valid upon restriction. A numerical condition involving the old modular step is not automatically unchanged when the new step is $\pm qd$.

There is a concrete obstruction to demanding long exact ordinary partitions from arbitrary proper modular parents. For any integer $m \geq 3$, take

$$N = 4m + 1, \quad P = \{4m\} \cup \{1, 3, \dots, 2m - 3\}.$$

This is a proper m -term modular progression of step 2, starting at $4m$. An ordinary progression contained in P and containing $4m$ has at most two terms. Otherwise its last three terms $a, b, 4m$ would satisfy $a \geq 1$ and $b = (a + 4m)/2 > 2m$, whereas every other member of P is at most $2m - 3$. Thus every exact ordinary-progression partition has a cell of size at most two, even as m grows and even if cells may use different steps. Removing the single outlier leaves an ordinary $(m - 1)$ -term progression, so this example establishes an obstruction to exact long partitions, not optimality of the square-root almost-cover bound. It concerns arbitrary modular input and does not refute a source theorem with stronger integer-geometric input.

This does not replace the current source-faithful Section 16 interface [119]. First, that interface requests an exact partition of the whole parent, whereas the theorem partitions only a retained subset. Adding singleton boxes covers the remainder but destroys a nontrivial uniform width bound. Second, the source's exceptional loss is on a base Q in a product $Q \times I$, whereas our loss is on product points. Deleting $Q \times \{y_0\}$ has product proportion $1/|I|$ but leaves no full original fiber untouched. Small product loss alone therefore does not control the base points whose full fibers survive. These two requirements need a separate construction or an explicitly changed interface followed through downstream uses. In a Section 16 product with k base axes and one last axis, the total dimension here is $D = k + 1$, so the denominator is $3(k + 1)$.

I.45 An explicit quadratic recurrence estimate

Source 32, delivered Section 2.

We give the exponential-sum argument in full, including the denominator and tail bounds. All logarithms are natural.

Lemma I.45.1 (Circle approximation). *For a real β and an integer $D \geq 1$, there are integers a, q with*

$$1 \leq q \leq D, \quad \gcd(a, q) = 1, \quad \left| \beta - \frac{a}{q} \right| \leq \frac{1}{qD}.$$

There is also, without the coprimality requirement, an integer $1 \leq q \leq D$ such that $\|q\beta\| \leq 1/(D + 1)$.

Proof. Consider $0, \beta, \dots, D\beta$ on the unit circle. If two coincide, their index difference gives distance zero. Otherwise some circular gap between neighboring points is at most $1/(D + 1)$. The absolute difference of their indices gives $1 \leq q_0 \leq D$ and an integer a_0 with $|q_0\beta - a_0| \leq 1/(D + 1)$. This also

covers a gap crossing zero. Reduce a_0/q_0 to a/q . Then $q \leq q_0 \leq D$ and

$$\left| \beta - \frac{a}{q} \right| \leq \frac{1}{q_0(D+1)} \leq \frac{1}{qD}.$$

The unreduced pair supplies the second assertion. $\square$

Lemma I.45.2 (Quadratic exponential sum). *Let $T \geq 2$ be an integer, let $\omega \in \mathbb{R}$, and put $\beta = 2\omega$. If*

$$\gcd(a, q) = 1, \quad 1 \leq q \leq T, \quad |\beta - a/q| \leq 1/(qT),$$

then

$$\left| \sum_{n=1}^T e(\omega n^2) \right|^2 \leq \frac{8T^2}{q} + 32T \log(2T). \quad (\text{I.45.1})$$

Proof. The case $q = 1$ follows from the trivial bound T^2 . Suppose $q \geq 2$ and put $b = \lfloor q/2 \rfloor \geq q/3$. Split $\{1, \dots, T-1\}$ into consecutive blocks of at most b terms. Their number is at most $T/b + 1 \leq 4T/q$.

For distinct indices u, v in the same block, $0 < |u - v| < q/2$. Coprimality and $q \leq T$ give

$$\|\beta(u - v)\| \geq \|a(u - v)/q\| - \frac{|u - v|}{qT} \geq \frac{1}{q} - \frac{1}{2T} \geq \frac{1}{2q}.$$

Thus the points βu belonging to one block have pairwise circle distance at least $1/(2q)$. Order their distances from zero as $d_1 \leq \dots \leq d_m$. An arc of length $2d$ contains at most $4qd + 1$ such points, so

$$d_j \geq \frac{j-1}{4q}.$$

Interpreting $1/0$ as infinity, it follows that

$$\begin{aligned} \sum_{u \text{ in block}} \min \left(T, \frac{1}{2\|\beta u\|} \right) &\leq T + \sum_{j=2}^m \frac{2q}{j-1} \\ &\leq T + 2q(1 + \log q). \end{aligned}$$

The last step uses $\sum_{j=1}^s 1/j \leq 1 + \log s$ for $s \geq 1$; the empty sum causes no problem. Summing over blocks gives

$$\sum_{u=1}^{T-1} \min \left(T, \frac{1}{2\|\beta u\|} \right) \leq \frac{4T^2}{q} + 8T(1 + \log q). \quad (\text{I.45.2})$$

For a geometric sum, $|1 - e(t)| = 2|\sin(\pi t)| \geq 4\|t\|$, and consequently

$$\left| \sum_{n=1}^K e(nt) \right| \leq \min \left(K, \frac{1}{2\|t\|} \right).$$

Expanding the squared quadratic sum and grouping pairs by their difference u now yields

$$\begin{aligned} \left| \sum_{n=1}^T e(\omega n^2) \right|^2 &\leq T + 2 \sum_{u=1}^{T-1} \left| \sum_{n=1}^{T-u} e(2\omega un) \right| \\ &\leq \frac{8T^2}{q} + 17T + 16T \log T \\ &\leq \frac{8T^2}{q} + 32T \log(2T). \end{aligned}$$

Here the constant phase $e(\omega u^2)$ has modulus one. The last inequality holds for $T \geq 2$, since $17 \leq 32 \log 2 + 16 \log T$. $\square$

Lemma I.45.3 (Uniform quadratic recurrence). *For every real θ and every integer $R \geq 2$, there is an integer p such that*

$$1 \leq p \leq 2^{10} R^5, \quad \|\theta p^2\| < 1/R. \quad (\text{I.45.3})$$

Proof. Set $T = 2^{10} R^5$ and $H = 4R^2$. Suppose for a contradiction that

$$\|\theta n^2\| \geq 1/R \quad (1 \leq n \leq T). \quad (\text{I.45.4})$$

On $\mathbb{R}/\mathbb{Z}$ consider the continuous triangle function

$$F(x) = \max\{1 - R\|x\|, 0\}.$$

Its Fourier coefficients are

$$c_0 = 1/R, \quad c_h = R \left(\frac{\sin(\pi h/R)}{\pi h} \right)^2 \geq 0 \quad (h \neq 0). \quad (\text{I.45.5})$$

For example, F is R times the circle convolution of the indicator of $[-1/(2R), 1/(2R)]$ with itself. Integrating that indicator against $e(-hx)$ and squaring gives (I.45.5). Also $c_h \leq R/(\pi^2 h^2)$, so the Fourier series is absolutely and uniformly convergent.

For completeness, this series indeed equals F . The Fejér kernel of order J is

$$K_J(x) = \frac{1}{J+1} \left| \sum_{v=0}^J e(vx) \right|^2.$$

It is nonnegative, its integral is one by termwise integration, and away from zero it is bounded by $1/((J+1)\sin^2(\pi x))$. Consequently convolution with K_J converges uniformly to any continuous periodic function: uniform continuity controls the near-zero integral, and the displayed bound controls the complement. Applied to F , these convolutions are also the Fejér averages of its absolutely convergent Fourier series, whose limit is the series itself. Hence $F(x) = \sum_h c_h e(hx)$, and at zero

$$\sum_{h \in \mathbb{Z}} c_h = F(0) = 1.$$

The high-frequency tail satisfies

$$\sum_{|h| > H} c_h \leq \frac{2R}{\pi^2} \sum_{h > H} h^{-2} \leq \frac{2R}{\pi^2 H} < \frac{1}{2R}. \quad (\text{I.45.6})$$

Write $S_h = \sum_{n=1}^T e(h\theta n^2)$. Assumption (I.45.4) says that $F(\theta n^2) = 0$ for all n in the sum. Interchanging a finite sum and the absolutely convergent Fourier series gives

$$0 = \frac{T}{R} + \sum_{h \neq 0} c_h S_h.$$

Using $|S_h| \leq T$ and (I.45.6), we obtain

$$\left| \sum_{0 < |h| \leq H} c_h S_h \right| > \frac{T}{2R}.$$

Since the coefficients in this sum have total at most one and $S_{-h} = \overline{S_h}$, some integer $1 \leq h \leq H$ satisfies

$$|S_h| > T/(2R). \quad (\text{I.45.7})$$

Apply Lemma I.45.1 with denominator budget T to $\beta = 2h\theta$. For the resulting relatively prime a, q , Lemma I.45.2 and (I.45.7) give

$$\frac{1}{4R^2} < \frac{8}{q} + \frac{32 \log(2T)}{T}. \quad (\text{I.45.8})$$

For $R \geq 2$, one has $\log R \leq R/2$: the difference $R/2 - \log R$ is positive at two and nondecreasing thereafter. Therefore

$$\log(2T) = 11 \log 2 + 5 \log R < 8R.$$

Indeed $11 \log 2 < 11 \leq 11R/2$. Since

$$T = 1024R^5 \geq 4096R^3 > 256R^2 \log(2T),$$

the second term on the right of (I.45.8) is less than $1/(8R^2)$. It follows that $q < 64R^2$.

Set $p = 2hq$. Then

$$1 \leq p < 512R^4 < T, \quad \|\beta q\| \leq 1/T,$$

and the exact identity $\theta p^2 = p(\beta q)$ gives

$$\|\theta p^2\| \leq p \|\beta q\| \leq \frac{p}{T} < \frac{1}{2R} < \frac{1}{R}.$$

This contradicts (I.45.4), since this integer p lies in its excluded interval. The contradiction proves (I.45.3). $\square$

[write] **6 October 2026; sources 28–33.** *Evidence, not part of the proof.* The intake's search over all $\theta = a/q$ with $q \leq 4R^2$ and 300 random θ for each $R = 2, \dots, 12$ found the least admissible p at most 53, far below $2^{10}R^5$: the bound is valid and very loose.

[write] **7 October 2026; sources 40–46.** Source 41 repeats this section verbatim (Section I.63) and extends the recurrence to every degree, $p \leq 2^{A_k} R^{B_k}$ (Theorem I.64.3).

Remark I.45.4. The bound $p < 512R^4$ in the last paragraph is obtained *under the contradiction hypothesis* (I.45.4). That hypothesis was used to force the large Fourier coefficient. It does not establish an unconditional R^4 recurrence theorem. The bound carried into every subsequent argument is $2^{10}R^5$.

I.46 Lossless phase partitions

Source 32, delivered Section 3.

The partitions below keep every point. A basic rounding fact is useful: a chain of length $s \geq L$ can be partitioned into chunks of lengths from L through $2L - 1$. Write $s = hL + b$ with $h \geq 1$ and $0 \leq b < L$, use $h - 1$ chunks of length L , and a last chunk of length $L + b$.

Lemma I.46.1 (Affine refinement). *Let m, L be integers with $1 \leq L \leq m$, and let $\ell, c \in \mathbb{R}$. The index interval $\{0, \dots, m-1\}$ has a partition into arithmetic progressions C of lengths $L \leq |C| \leq 2L-1$ and complex units z_C for which*

$$|e(-(\ell j + c)) - z_C| \leq \frac{2\pi L^2}{m} \quad (j \in C). \quad (\text{I.46.1})$$

Proof. Put $D = \lfloor m/L \rfloor \geq 1$. By the circle part of Lemma I.45.1, choose $1 \leq q \leq D$ with $\|q\ell\| \leq 1/(D+1)$. Split the interval into its residue chains modulo q . Every chain has at least $\lfloor m/q \rfloor \geq L$ points. Apply the chunking rule above to each chain.

Write $q\ell = v + \epsilon$, where $v \in \mathbb{Z}$ and $|\epsilon| \leq 1/(D+1)$. For a chunk $j = j_0 + tq$, $0 \leq t < s$, set

$$z_C = e(-(\ell j_0 + c))e(-\epsilon(s-1)/2). \quad (\text{I.46.2})$$

The integer vt does not affect the exponential. The chord-arc inequality gives

$$\begin{aligned} |e(-(\ell j + c)) - z_C| &\leq 2\pi|\epsilon| \left| t - \frac{s-1}{2} \right| \\ &\leq \pi|\epsilon|(s-1) \leq \frac{2\pi(L-1)}{D+1} < \frac{2\pi L^2}{m} \end{aligned}$$

when $L > 1$, using $D+1 > m/L$. When $L = 1$, every chunk is a singleton and the error is zero. This proves the stated weak bound in all cases. $\square$

The reduction before taking the midpoint in (I.46.2) matters. If $q\ell = 1$ and $s = 2$, both sampled phase values agree, while the midpoint obtained from the unreduced increment contributes $e(-1/2) = -1$. The reduced increment is zero and its midpoint has the correct phase. Integer shifts are harmless at integer parameters but need not be harmless at half-integer midpoints.

Lemma I.46.2 (Quadratic refinement). *Let $L \geq 2$ and $n \geq 2^{116}L^{38}$ be integers. Let $P = \{u + jd : 0 \leq j < n\}$ be a proper modular progression in G and let $\phi(x) = ax^2 + bx + c$. There is a partition of P into proper modular subprogressions C and complex units z_C such that*

$$L \leq |C| \leq 2L-1, \quad |e_N(-\phi(x)) - z_C| < \frac{1}{4L} \quad (x \in C). \quad (\text{I.46.3})$$

Proof. Choose integer representatives u, d, a, b, c . Define the real polynomial

$$p_0(j) = \theta j^2 + \beta j + \gamma, \quad \theta = \frac{ad^2}{N}, \quad \beta = \frac{2aud + bd}{N}, \quad \gamma = \frac{au^2 + bu + c}{N}.$$

Then $e_N(-\phi(u + jd)) = e(-p_0(j))$ for every integer j . The identity remains true when $u + jd$ wraps modulo N : replacing an integer representative changes the normalized polynomial value by an integer. No division inside G is involved, and neither d nor 2 needs to be a unit.

Set

$$H = 64L^3 = 2^6L^3, \quad R = 256LH^2 = 2^{20}L^7, \quad T = 2^{10}R^5 = 2^{110}L^{35}. \quad (\text{I.46.4})$$

Lemma I.45.3 supplies $1 \leq p \leq T$ with $\|\theta p^2\| < 1/R$. Since $n \geq TH$, each residue chain modulo p in the parent index interval has at least $\lfloor n/p \rfloor \geq H$ points. Chunk each chain into pieces of lengths H through $2H-1$.

On one piece $j = j_0 + pt$, $0 \leq t < m$, $H \leq m \leq 2H - 1$, expansion gives

$$p_0(j_0 + pt) = \theta p^2 t^2 + \ell t + c_0, \quad \ell = 2\theta j_0 p + \beta p, \quad c_0 = p_0(j_0).$$

Choose $v \in \mathbb{Z}$ so that $\epsilon = \theta p^2 - v$ has $|\epsilon| < 1/R$. Subtracting vt^2 leaves the exponential unchanged at integer t . Since $t < 2H$, the curvature error satisfies

$$|e(-p_0(j_0 + pt)) - e(-(\ell t + c_0))| \leq 2\pi|\epsilon|t^2 < \frac{8\pi H^2}{R} = \frac{\pi}{32L}. \quad (\text{I.46.5})$$

Apply Lemma I.46.1 to the affine phase on this t -interval with target length L . Its extra error is at most

$$\frac{2\pi L^2}{m} \leq \frac{2\pi L^2}{H} = \frac{\pi}{32L}. \quad (\text{I.46.6})$$

Combining (I.46.5) and (I.46.6) gives error less than $\pi/(16L) < 1/(4L)$.

Each final cell has a parameterization $t = t_0 + qs$ within a first-stage piece, hence $j = j_0 + pt_0 + pqs$ in the original interval. Its image in G is a modular progression. The parent index map is injective, so its restriction to each final cell is injective and preserves cardinality. Distinct cells stay disjoint, and all parent indices are retained. Finally,

$$TH = 2^{116} L^{38},$$

which is exactly the hypothesis. □

I.47 A second moment over every cyclic modulus

Source 32, delivered Section 4.

Over a prime modulus, pairs of distinct progression positions often have a particularly simple distribution. Over an arbitrary cyclic modulus their distribution is instead controlled by additive subgroup cosets. The resulting inequality is enough here.

Lemma I.47.1 (Proper progression from variance). *Let $0 < \delta < 1$, let $L \geq 2$ be an integer, and suppose $N \geq L^3$. There is a proper L -term modular progression Q with*

$$\text{dens}_Q(A) := \frac{|A \cap Q|}{|Q|} \geq \delta + \frac{1 - \delta}{2L}. \quad (\text{I.47.1})$$

Proof. Choose a, d independently and uniformly in G , including $d = 0$, and put

$$X(a, d) = \frac{1}{L} \sum_{j=0}^{L-1} \mathbf{1}_A(a + jd).$$

Repeated positions are counted with their index multiplicity. For every fixed d , averaging over a gives $\mathbb{E}_a X = \delta$.

For distinct j, k , let $H = (k - j)G$. The change of variables $(a, d) \mapsto (x, d) = (a + jd, d)$ is a bijection of G^2 . The element $(k - j)d$ is uniform in H , independently of x . Therefore

$$\mathbb{E}[\mathbf{1}_A(a + jd)\mathbf{1}_A(a + kd)] = \sum_{C \in G/H} \frac{|H|}{N} \left(\frac{|A \cap C|}{|H|} \right)^2 \geq \delta^2. \quad (\text{I.47.2})$$

Indeed the coset densities have weighted mean δ , and their weighted squared deviations from that mean sum to a nonnegative number. Thus all off-diagonal covariances are nonnegative, and the L diagonal terms yield

$$\text{Var}_{a,d} X \geq \frac{\delta(1-\delta)}{L}. \quad (\text{I.47.3})$$

A direction d is bad precisely when the L -term parameterization is not injective. This means $hd = 0$ for some $1 \leq h < L$. The kernel of multiplication by h on G has $\gcd(h, N) \leq h$ elements: if $g = \gcd(h, N)$, then $N \mid hd$ is equivalent to $N/g \mid d$. Hence the proportion η of bad directions obeys

$$\eta \leq \frac{1}{N} \sum_{h=1}^{L-1} h = \frac{L(L-1)}{2N} < \frac{1}{2L}. \quad (\text{I.47.4})$$

At least one proper direction exists, since $d = 1$ works for $L \leq N$.

Let h_* be the largest value of $X(a, d) - \delta$ over proper directions and all a . For a proper direction, $0 \leq X \leq \delta + h_*$ implies $X^2 \leq (\delta + h_*)X$, and so $\text{Var}_a X \leq \delta h_*$. For any bad direction, $0 \leq X \leq 1$ similarly gives $\text{Var}_a X \leq \delta(1-\delta)$. Since every conditional mean is δ , averaging these variances and using (I.47.3) gives

$$\frac{\delta(1-\delta)}{L} \leq (1-\eta)\delta h_* + \eta\delta(1-\delta).$$

Dividing by $\delta > 0$ and rearranging,

$$h_* \geq (1-\delta) \frac{1/L - \eta}{1-\eta} \geq \frac{1-\delta}{2L},$$

where (I.47.4) and $1-\eta \leq 1$ justify the last step. A maximizing proper progression proves the claim. $\square$

[write] **6 October 2026; sources 28–33.** Source 26's Lemma I.42.1 is the same second-moment argument over every cyclic modulus (the sibling reports share this section); source 42's placement records a version needing only $N \geq 2(L-1)^2 + 1$ (for its write).

I.48 Assembly and every scale

Source 32, delivered Section 5.

Proof of Theorem F.37.1. Let

$$b = 1 - \delta, \quad \tau = |U|/N, \quad r = \min_i |P_i|, \quad S = N/M, \quad \mu = \frac{1}{N} \sum_{x \in U} |f(x)|.$$

Because the parents are nonempty and disjoint, $M \leq N$ and $S \geq 1$. The cover hypothesis implies $\delta \leq \tau \leq 1$. Separating the δN points of A from the $(\tau - \delta)N$ other points of U , we obtain

$$\mu = \delta(1 + \tau - 2\delta), \quad 0 < \alpha \leq \mu \leq 2\delta b \leq 2\tau b \leq 2b. \quad (\text{I.48.1})$$

Here $\alpha \leq \mu$ follows from the correlation and the triangle inequality. The positive correlation also rules out $\delta = 0$ and $\delta = 1$, so $0 < \delta < 1$, $b > 0$, and A is nonempty. Comparability gives

$$\tau N = \sum_i |P_i| \leq 2Mr, \quad r \geq \frac{\tau S}{2}. \quad (\text{I.48.2})$$

Set

$$x = \frac{S^{1/39}}{17}, \quad L = \lceil x \rceil. \quad (\text{I.48.3})$$

If $L = 1$, choose a singleton in A . It has length at least x and density one. By (I.48.1), $\alpha/3 \leq 2\delta b/3 < b$, so its density is strictly greater than $\delta + \alpha/3$.

Now suppose $L \geq 2$. Then $x > 1$ and $L < x + 1 < 2x$, so

$$S > (17L/2)^{39} > L^3, \quad N \geq S > L^3. \quad (\text{I.48.4})$$

If $\alpha \leq 3b/(2L)$, Lemma I.47.1 gives a proper progression of length L with density at least

$$\delta + \frac{b}{2L} \geq \delta + \frac{\alpha}{3}.$$

Its length is at least x , as required.

It remains to consider $\alpha > 3b/(2L)$. Since $\alpha \leq 2\tau b$, this implies $\tau > 3/(4L)$. Equations (I.48.2) and (I.48.4) give

$$r > \frac{3S}{8L} > \frac{3 \cdot 17^{39}}{2^{42}} L^{38} > 2^{116} L^{38}. \quad (\text{I.48.5})$$

The last numerical inequality follows entirely with integers. Bernoulli's inequality gives

$$17^{39} = 2^{156} (1 + 1/16)^{39} \geq 2^{156} (1 + 39/16) = \frac{55}{16} 2^{156},$$

and $3 \cdot 55/16 > 4$, hence $3 \cdot 17^{39} > 2^{158}$.

Apply Lemma I.46.2 to every parent at this common target length L . For each resulting cell C , let $i(C)$ denote its parent, let z_C be its approximating unit, and set $F_C = \sum_{x \in C} f(x) \in \mathbb{R}$. The total phase error is

$$\begin{aligned} E &:= \sum_C \sum_{x \in C} |f(x)| |e_N(-\phi_{i(C)}(x)) - z_C| \\ &< \frac{\mu N}{4L} \leq \frac{bN}{2L} < \frac{\alpha N}{3}. \end{aligned} \quad (\text{I.48.6})$$

For each parent, replacing its phase on each cell by z_C and then applying the triangle inequality bounds its correlation by the sum of its $|F_C|$ plus its phase error. Summing over the parents and using (F.37.2) yields

$$\sum_C |F_C| \geq \alpha N - E > 2\alpha N/3. \quad (\text{I.48.7})$$

The signed total has a favorable sign, because $A \subseteq U$:

$$\sum_C F_C = |A| - \delta|U| = \delta(N - |U|) \geq 0. \quad (\text{I.48.8})$$

The sum of the positive F_C is half the sum of (I.48.7) and (I.48.8), and is therefore greater than $\alpha N/3$. These positive cells together have at most N points. If each had $F_C/|C| \leq \alpha/3$, their positive total would be at most $\alpha N/3$, a contradiction. Thus some cell Q satisfies

$$\frac{|A \cap Q|}{|Q|} = \delta + \frac{F_Q}{|Q|} > \delta + \frac{\alpha}{3}.$$

This cell is proper and has length at least $L \geq x$. In all branches the progression is nonempty, which also proves the maximum with one in (F.37.3). $\square$

The large-correlation branch finds its progression inside a parent. The variance branch searches all modular directions in G , so the theorem as a whole does not promise parent containment. The word “threshold free” refers to the theorem’s hypotheses; the internal phase-partition lemma has its explicit threshold, and the preceding dichotomy supplies it exactly where needed.

I.49 Source 37’s chapter: missing intervals at Lemmas 5.2 and 5.5 [write]

Added in the seventh write, 6 October 2026.

The next ten sections are source 37’s Sections 2–9 and Appendices A–B, on the missing-interval step of Gowers’s Section 5: a set, or a measure, that avoids an arc around zero must have a large low Fourier coefficient. *Logically the chapter belongs before the partition material of Section 5* (Lemma 5.2 precedes the polynomial recurrence of Lemma 5.5 and the partitions of Corollary 5.6); it is printed here, at the end of Part I, because no number printed earlier may move. No earlier source of this report treats Lemma 5.2. Source 37’s introduction is Section F.42; its verification, questions and source audit are in Part V.

What is proved, in one paragraph. For $\pi < (R + 1)a < 2\pi$ the least possible value of $\max_{1 \leq r \leq R} |\hat{\mu}(r)|$ over probability measures on the circle avoiding the open arc of radius a is $\rho_R(a) = 1/(R - \cot(a/2) \tan((R + 1)a/2))$, attained by a unique measure with $R + 1$ atoms, all of whose first R coefficients equal $-\rho_R(a)$; with coefficients at most ε the least mass of the arc is $\max\{0, (\rho - \varepsilon)/(1 + \rho)\}$ (Theorem I.50.1), with an explicit polynomial certificate (Section I.51), a moment criterion (Section I.52), the critical radius (Theorem I.54.1) and stability (Theorems I.56.1, I.56.3). The universal consequence (Theorem I.55.1) gives a replacement for Lemma 5.2 on $\mathbb{Z}/N\mathbb{Z}$ with frequency at most $\lfloor N/M \rfloor - 1$ (not $(N/M)^2$) and coefficient $|A|M/(N - M)$ (not $|A|M/(4N)$), for arbitrary nonnegative weights and without the parity of M (Corollary I.55.2); with the explicit Weyl input recorded in the formal `lemma_5_3` it gives the recurrence exponent $b_k = (1 - 1/(4k))/(2^{k-1} + 1)$ ($7/24$ at $k = 2$, against Lemma 5.5’s $1/16$; Corollary I.57.3) and a short-step version (Corollary I.57.4). The general moment criterion is classical (Yang and Xie are credited); the Weyl inequality is imported.

Relation to the formal statements (checked at HEAD 89c1cbc38). `lemma_5_2` (Section05.lean, line 141) and `lemma_5_5` (line 181) have the exact companions `lemma_5_2_holds` (Proofs05MissingInterval.lean, line 220) and `lemma_5_5_holds` (Proofs05PolynomialPartition.lean, line 1678). Corollary I.55.2 implies the conclusion of `lemma_5_2`, and Corollary I.57.3 the bound of `lemma_5_5` for $t \geq \text{weylThreshold } k$ (proved at the write; notes there). Nothing in this chapter is formalized; the status of both statements is unchanged.

I.50 The exact gap–moment theorem

Source 37, delivered Section 2.

I.50.1 Conventions

Let $\mathbb{T}_{2\pi} = \mathbb{R}/(2\pi\mathbb{Z})$, with angular distance $|\theta|_o$ from zero. Write

$$I_a = \{\theta \in \mathbb{T}_{2\pi} : |\theta|_o < a\}, \quad K_a = \mathbb{T}_{2\pi} \setminus I_a, \quad 0 < a \leq \pi.$$

The arc is *open*; its endpoints belong to K_a . All measures are Borel probability measures unless a different total mass is specified; $\mathcal{P}(K_a)$ denotes the probability measures supported on K_a . Define

$$\hat{\mu}(r) = \int_{\mathbb{T}_{2\pi}} e^{-ir\theta} d\mu(\theta), \quad E_R(\mu) = \max_{1 \leq r \leq R} |\hat{\mu}(r)|. \quad (\text{I.50.1})$$

No absolute continuity, distinctness of sample points, rationality of angles, or primality of an underlying modulus is assumed.

For an integer $R \geq 1$, set $m = R + 1$. In the *strict first window*

$$\frac{\pi}{m} < a < \frac{2\pi}{m}, \quad (\text{I.50.2})$$

put

$$S_R(a) = m - \cot(a/2) \tan(ma/2), \quad (\text{I.50.3})$$

$$d_R(a) = S_R(a)^{-1}, \quad \rho_R(a) = \frac{d_R(a)}{1 - d_R(a)} = \frac{1}{R - \cot(a/2) \tan(ma/2)}. \quad (\text{I.50.4})$$

Since $\tan(ma/2) < 0$, one has $S_R(a) > m$ and $0 < \rho_R(a) < 1/R$.

Theorem I.50.1 (Exact first-window tradeoff). *Suppose (I.50.2) holds, and abbreviate $d = d_R(a)$ and $\rho = \rho_R(a)$.*

(i) *For every probability measure μ ,*

$$\mu(I_a) \geq d - (1 - d)E_R(\mu). \quad (\text{I.50.5})$$

If $\mu(I_a) = 0$, then for some $1 \leq r \leq R$,

$$\operatorname{Re} \hat{\mu}(r) \leq -\rho. \quad (\text{I.50.6})$$

(ii) *There is a unique probability measure $\mu_{R,a}$ supported on K_a with $E_R(\mu_{R,a}) = \rho$. It satisfies*

$$\hat{\mu}_{R,a}(r) = -\rho \quad (1 \leq r \leq R), \quad (\text{I.50.7})$$

and has exactly $R + 1$ support points with strictly positive masses. Consequently,

$$\inf_{\mu \in \mathcal{P}(K_a)} E_R(\mu) = \inf_{\mu \in \mathcal{P}(K_a)} \max_{1 \leq r \leq R} (-\operatorname{Re} \hat{\mu}(r)) = \rho. \quad (\text{I.50.8})$$

(iii) *For every $\varepsilon \geq 0$,*

$$\inf_{E_R(\mu) \leq \varepsilon} \mu(I_a) = \max \left\{ 0, \frac{\rho - \varepsilon}{1 + \rho} \right\}. \quad (\text{I.50.9})$$

For $0 \leq \varepsilon \leq \rho$, the unique minimizing measure is

$$\mu_\varepsilon = s\delta_0 + (1 - s)\mu_{R,a}, \quad s = \frac{\rho - \varepsilon}{1 + \rho}. \quad (\text{I.50.10})$$

[write] **7 October 2026; sources 60–69.** *Unequal budgets* (source 67, Section I.75): when the R coefficient budgets differ, the least mass is $-1/((1 - \cos a) \mathbf{e}^T L(\boldsymbol{\varepsilon})^{-1} \mathbf{e})$ on an explicit neighbourhood of every interior point of this theorem's equal-budget branch, and this theorem's certificate stays optimal exactly on a polytope of local dimension $\lfloor R/2 \rfloor + 1$ (its Theorems I.76.2–I.76.5).

The equality case is part of the assertion: the bound is not simply the best one obtained from a chosen family of test polynomials. It is optimal over all probability measures. The proof occupies Sections I.51–I.53. Endpoint values and the critical geometry are treated separately in Section I.54.

I.50.2 Description of the extremizer

The support has an explicit algebraic description. Define

$$w_j = \frac{1 - \cos((j - (R - 1)/2)a) / \cos(ma/2)}{S_R(a)}, \quad 0 \leq j \leq R - 1, \quad (\text{I.50.11})$$

and

$$P_{R,a}(z) = \sum_{j=0}^{R-1} w_j z^j. \quad (\text{I.50.12})$$

Then $\text{supp } \mu_{R,a}$ consists of $-a$, a , and the arguments of the $R - 1$ zeros of $P_{R,a}$. All of those zeros are simple and lie on the unit circle strictly inside the complementary arc. These root facts are proved, not assumed from numerics. The masses are the unique solution of

$$\sum_{j=1}^{R+1} u_j = 1, \quad \sum_{j=1}^{R+1} u_j e^{-ir\theta_j} = -\rho \quad (1 \leq r \leq R). \quad (\text{I.50.13})$$

The matrix of this system is Vandermonde and hence nonsingular. The proof also establishes $u_j > 0$ and reflection symmetry. Thus root finding followed by a small linear solve gives a numerical realization of the exact measure.

I.51 A positive-coefficient polynomial certificate

Source 37, delivered Section 3.

I.51.1 A path recurrence

Let T_R be the $R \times R$ real symmetric matrix with $1/2$ on each adjacent off-diagonal and zero elsewhere. Write $c = \cos a$, $S = S_R(a)$, and

$$\lambda = \frac{1 - c}{S} = (1 - c)d. \quad (\text{I.51.1})$$

Lemma I.51.1 (The coefficient vector). *The vector in (I.50.11) has positive entries, is symmetric under $j \mapsto R - 1 - j$, and satisfies*

$$\sum_{j=0}^{R-1} w_j = 1, \quad (T_R - cI)w = \lambda \mathbf{1}. \quad (\text{I.51.2})$$

Proof. In the first window $\cos(ma/2) < 0$. Moreover, $|(j - (R - 1)/2)a| < ma/2 < \pi$, so its cosine is strictly larger than $\cos(ma/2)$. Dividing by the negative denominator proves $w_j > 0$. Symmetry is immediate. The finite cosine sum gives

$$\sum_{j=0}^{R-1} \cos((j - (R - 1)/2)a) = \frac{\sin(Ra/2)}{\sin(a/2)}.$$

Since $R = m - 1$,

$$\frac{\sin(Ra/2)}{\sin(a/2) \cos(ma/2)} = \cot(a/2) \tan(ma/2) - 1.$$

The sum of the numerators in (I.50.11) is therefore S . For the recurrence, extend the displayed formula to $j = -1, R$; it gives $w_{-1} = w_R = 0$. The cosine addition formula then yields

$$\frac{w_{j-1} + w_{j+1}}{2} - cw_j = \frac{1-c}{S} = \lambda \quad (0 \leq j \leq R-1).$$

This includes both boundary coordinates of the path matrix. $\square$

I.51.2 The crucial sign computation

Define a real even trigonometric polynomial of degree R by

$$Q(\theta) = (\cos \theta - c)|P_{R,a}(e^{i\theta})|^2 = q_0 + 2 \sum_{k=1}^R q_k \cos(k\theta). \quad (\text{I.51.3})$$

The sign of Q is already correct for detecting the gap. The useful extra fact is positivity of *every* nonconstant Fourier coefficient.

Lemma I.51.2 (Exact Fourier coefficients). *One has $q_0 = \lambda$ and, for $1 \leq k \leq R$,*

$$q_k = \lambda \sum_{j=0}^{R-k-1} w_j + \frac{1}{2} w_{R-k} w_{R-1} > 0. \quad (\text{I.51.4})$$

An empty sum is zero. In particular $q_R = w_0 w_{R-1}/2 > 0$.

Proof. Extend w_j by zero outside $0 \leq j < R$ and put $A_k = \sum_j w_j w_{j+k}$, so $A_{-k} = A_k$. Multiplication in (I.51.3) gives

$$q_k = \frac{1}{2}(A_{k-1} + A_{k+1}) - cA_k = \sum_{j=0}^{R-1} w_j \left(\frac{w_{j+k-1} + w_{j+k+1}}{2} - cw_{j+k} \right).$$

For $k = 0$, every bracket equals λ , so $q_0 = \lambda$. For $1 \leq k \leq R$, the brackets with $0 \leq j < R - k$ equal λ ; the bracket at $j = R - k$ equals $w_{R-1}/2$; all later brackets vanish. This proves (I.51.4). Positivity follows from Lemma I.51.1 and $\lambda > 0$. $\square$

Normalize by

$$F(\theta) = \frac{Q(\theta)}{1-c} = d + 2 \sum_{k=1}^R f_k \cos(k\theta), \quad f_k = \frac{q_k}{1-c}. \quad (\text{I.51.5})$$

Because $P(1) = 1$, we have $F(0) = 1$. Thus

$$f_k > 0, \quad 2 \sum_{k=1}^R f_k = 1 - d. \quad (\text{I.51.6})$$

On K_a , $F \leq 0$. On the whole circle, positivity of the Fourier coefficients gives $F(\theta) \leq F(0) = 1$. Consequently

$$F \leq \mathbf{1}_{I_a}. \quad (\text{I.51.7})$$

The value at either endpoint is zero, so the open-arc convention causes no error.

Proposition I.51.3 (Mass and one-sided spectral inequalities). *For any probability measure μ ,*

$$\mu(I_a) \geq d + 2 \sum_{k=1}^R f_k \operatorname{Re} \hat{\mu}(k). \quad (\text{I.51.8})$$

In particular (I.50.5) holds. If $\mu(I_a) = 0$, at least one $\operatorname{Re} \hat{\mu}(k)$ is at most $-\rho$.

Proof. Integrate (I.51.7). Each real part is at least $-E_R(\mu)$; (I.51.6) gives (I.50.5). If all real parts were strictly greater than $-\rho$, then the integral of F would be strictly greater than $d - (1 - d)\rho = 0$, contradicting $F \leq 0$ on the support. $\square$

Remark I.51.4 (Unequal moment information). The unsimplified inequality (I.51.8) is useful when frequencies have different error bounds. From $|\hat{\mu}(k)| \leq \varepsilon_k$ it gives $\mu(I_a) \geq d - 2 \sum f_k \varepsilon_k$. We do not claim that this last anisotropic bound is optimal for an arbitrary vector of caps. Exact optimality is proved here for equal caps.

I.52 A supported trigonometric moment criterion

Source 37, delivered Section 4.

Attainment requires more than a dual certificate. We give the relevant finite-dimensional moment result and its proof, so that the construction of the extremal measure does not rest on an unexplained semidefinite program. This criterion is classical in substance; compare [152].

I.52.1 Positive trigonometric polynomials on one arc

Lemma I.52.1 (Arc positivity representation). *Let $0 < a < \pi$, $c = \cos a$, and let H be a real-valued trigonometric polynomial of degree at most R . If $H \geq 0$ on K_a , then there are complex polynomials A, B with $\deg A \leq R$ and $\deg B \leq R - 1$ such that*

$$H(\theta) = |A(e^{i\theta})|^2 + (c - \cos \theta) |B(e^{i\theta})|^2. \quad (\text{I.52.1})$$

Proof. Set $t = \cot(\theta/2)$ and $b = \cot(a/2) > 0$. The closed arc K_a maps to $[-b, b]$. Because $e^{i\theta} = (t + i)/(t - i)$,

$$p(t) = (1 + t^2)^R H(\theta(t))$$

is a real polynomial of degree at most $2R$, nonnegative on $[-b, b]$. The elementary interval factorization proved in Appendix I.58 gives

$$p(t) = U(t)^2 + (b^2 - t^2)V(t)^2, \quad \deg U \leq R, \quad \deg V \leq R - 1.$$

Now $U(t)/(t - i)^R$ and $V(t)/(t - i)^{R-1}$ are polynomials in $z = (t + i)/(t - i)$ of the indicated degrees. Also

$$\frac{b^2 - t^2}{1 + t^2} = \frac{c - \cos \theta}{1 - c}.$$

Divide the factorization by $(1 + t^2)^R$ and absorb the positive factor $(1 - c)^{-1/2}$ into the second polynomial. The identity extends to $\theta = 0$ by continuity. $\square$

Theorem I.52.2 (One-arc moment criterion). *Let complex numbers $h_{-R}, \dots, h_R$ satisfy $h_0 = 1$ and $h_{-j} = \overline{h_j}$. They are the moments $h_j = \hat{\mu}(j)$ of a probability measure supported on K_a if and only if both matrices*

$$M = (h_{j-k})_{0 \leq j, k \leq R}, \quad (\text{I.52.2})$$

$$L = \left(ch_{j-k} - \frac{h_{j-k+1} + h_{j-k-1}}{2} \right)_{0 \leq j, k < R} \quad (\text{I.52.3})$$

are positive semidefinite.

Proof. Necessity follows by integrating respectively $|A(e^{i\theta})|^2$ and $(c - \cos \theta)|B(e^{i\theta})|^2$ against μ . For sufficiency, let $\mathcal{L}$ be the real-linear functional on real trigonometric polynomials of degree at most R determined by these moments. Its value on a constant is that constant. The two matrix inequalities imply that it is nonnegative on both terms in (I.52.1), and hence on every such polynomial nonnegative on K_a .

Consider the compact curve

$$v(\theta) = (\cos \theta, -\sin \theta, \dots, \cos R\theta, -\sin R\theta) \in \mathbb{R}^{2R}, \quad \theta \in K_a.$$

Its convex hull is compact: by the finite-dimensional Carathéodory argument, every convex combination can be reduced to at most $2R + 1$ points by removing an affine dependence; compactness then follows from the compactness of the parameter spaces for those combinations. If the vector $(\operatorname{Re} h_1, \Im h_1, \dots, \operatorname{Re} h_R, \Im h_R)$ lay outside this hull, a strict separating affine functional would give a real trigonometric polynomial H nonnegative on K_a but with $\mathcal{L}(H) < 0$. This is impossible. A convex combination representing that vector gives the desired probability measure. $\square$

I.52.2 The needed spectral inequality

Lemma I.52.3 (Zero-sum path bound). *For $v \in \mathbb{C}^R$ with $\sum_j v_j = 0$,*

$$v^* T_R v \leq \cos \left(\frac{2\pi}{R+1} \right) \|v\|_2^2. \quad (\text{I.52.4})$$

Proof. Append one zero coordinate to obtain a vector $y \in \mathbb{C}^{R+1}$ of mean zero. Let C_m act on cyclic coordinates by $(C_m y)_j = (y_{j-1} + y_{j+1})/2$. Then $y^* C_m y = v^* T_R v$, since every term involving the appended coordinate is zero. The cyclic Fourier vectors have eigenvalues $\cos(2\pi\ell/m)$, $0 \leq \ell < m$. The mean-zero condition removes $\ell = 0$, so the largest remaining eigenvalue is $\cos(2\pi/m)$. This proves the inequality. For $R = 1$ the only zero-sum vector is zero, and the same conclusion holds. $\square$

The use of the second cyclic eigenvalue is essential. On the full path space the largest eigenvalue is $\cos(\pi/m)$, which exceeds $\cos a$ in the first window. The rank-one correction below removes precisely that bad one-dimensional direction.

I.53 Attainment, support, and equality

Source 37, delivered Section 5.

I.53.1 Constructing the negative equimoment measure

We prescribe

$$h_0 = 1, \quad h_j = -\rho \quad (1 \leq |j| \leq R). \quad (\text{I.53.1})$$

For these moments the Toeplitz matrix is

$$M = (1 + \rho)I_m - \rho J_m, \quad (\text{I.53.2})$$

where $J_m = \mathbf{1}\mathbf{1}^*$. Its eigenvalues are $1 + \rho$ on $\mathbf{1}^\perp$ and $1 - R\rho$ on $\mathbf{1}$. Both are positive in the strict first window.

The localizing matrix has the explicit form

$$\begin{aligned} L &= (1 + \rho)(cI_R - T_R) + \rho(1 - c)J_R \\ &= (1 + \rho)(cI_R - T_R + \lambda J_R), \end{aligned} \quad (\text{I.53.3})$$

since $\lambda = (1 - c)\rho/(1 + \rho)$.

Lemma I.53.1 (The rank-one correction is positive). *The matrix $K = cI_R - T_R + \lambda J_R$ is positive semidefinite, has kernel $\mathbb{C}w$, and has rank $R - 1$.*

Proof. Equation (I.51.2) and $\sum w_j = 1$ give $Kw = 0$. Every $x \in \mathbb{C}^R$ can be written uniquely as

$$x = u + \left(\sum_j x_j \right) w, \quad \sum_j u_j = 0.$$

Since K is Hermitian and $Kw = 0$,

$$x^* K x = u^* K u = u^* (cI_R - T_R) u.$$

The first-window upper inequality gives $c > \cos(2\pi/m)$, because $0 < a < 2\pi/m \leq \pi$. Lemma I.52.3 shows that the last quadratic form is positive for every nonzero u . This proves positivity, the stated kernel, and the rank. $\square$

Theorem I.52.2 now produces a probability measure μ on K_a with the prescribed moments. This establishes attainment of the lower bound, but there is considerably more structure.

I.53.2 Why all roots lie on the circle

Proposition I.53.2 (Support and uniqueness). *The prescribed measure exists uniquely. It has the support and positive weights described in (I.50.12)–(I.50.13). In particular, $P_{R,a}$ has $R - 1$ distinct unit-circle zeros in the relative interior of K_a , and neither endpoint is a zero.*

Proof. For any representing measure on K_a ,

$$\int (c - \cos \theta) |P_{R,a}(e^{i\theta})|^2 d\mu = w^* L w = 0.$$

The integrand is nonnegative. Thus the support is contained in the two endpoints and the unit-circle zeros of $P_{R,a}$. Since P has degree $R - 1$ and nonzero leading coefficient, there are at most $R + 1$ distinct possible support points.

On the other hand M in (I.53.2) is positive definite and has rank $R + 1$. A measure supported at s points has a moment matrix that is a sum of s rank-one matrices and therefore has rank at most s . Hence at least $R + 1$ support points are necessary. Every inequality is therefore an equality: both endpoints occur; all $R - 1$ roots occur, are distinct, lie on the unit circle, and are not endpoints; and every mass is positive.

For this fixed support, the equations for $h_0, \dots, h_R$ form a square Vandermonde system in the distinct numbers $e^{-i\theta_j}$. Its determinant is nonzero, so the weights are unique. The support was determined by P , so the measure is unique. Reflecting the measure changes each moment to its complex conjugate and preserves the support; the moments are real, so uniqueness also proves reflection symmetry. $\square$

This argument is also a proof of the root-location assertion. Numerical root plots are not being used as a substitute for a unit-circle theorem.

I.53.3 Equality in the minimax and noisy problems

If $\mu \in \mathcal{P}(K_a)$ and $E_R(\mu) = \rho$, the inequalities in Proposition I.51.3 are all equalities. Since every $f_k > 0$ and $\operatorname{Re} \hat{\mu}(k) \geq -\rho$, each real part is $-\rho$. Together with $|\hat{\mu}(k)| \leq \rho$, this forces $\hat{\mu}(k) = -\rho$. Proposition I.53.2 proves the uniqueness claimed in Theorem I.50.1(ii). The same argument proves the one-sided minimax identity.

For $0 \leq \varepsilon \leq \rho$, the measure in (I.50.10) has

$$\hat{\mu}_\varepsilon(k) = s - (1 - s)\rho = -\varepsilon, \quad \mu_\varepsilon(I_a) = s.$$

It therefore attains the lower bound in (I.50.5). If $\varepsilon \geq \rho$, $\mu_{R,a}$ gives gap mass zero. This proves the value in (I.50.9).

For completeness, suppose $0 \leq \varepsilon \leq \rho$ and equality is attained. Every nonnegative term in

$$\mu(I_a) - [d - (1 - d)\varepsilon] = \int (\mathbf{1}_{I_a} - F) d\mu + 2 \sum_{k=1}^R f_k (\operatorname{Re} \hat{\mu}(k) + \varepsilon) \quad (\text{I.53.4})$$

must vanish. Thus all moments equal $-\varepsilon$. Inside I_a , equality $F = 1$ is possible only at zero: the positive coefficient f_1 already forces $\cos \theta = 1$. Outside, equality $F = 0$ occurs precisely at the $R + 1$ support points from Proposition I.53.2. Consequently any equality measure is supported on the fixed $R + 2$ -point set

$$Z = \{0\} \cup \operatorname{supp} \mu_{R,a}. \quad (\text{I.53.5})$$

The matrix with columns $e^{-ik\theta_j}$, $-R \leq k \leq R$, has full column rank: multiplying each column by $e^{-iR\theta_j}$ turns it into a Vandermonde matrix with powers $0, \dots, 2R$, and $R + 2 \leq 2R + 1$. Thus the moments determine the weights uniquely. The measure (I.50.10) is that unique measure. This completes the proof of Theorem I.50.1.

I.54 Endpoints and the first spectral transition

Source 37, delivered Section 6.

I.54.1 Zero moments and the critical gap

Theorem I.54.1 (The critical radius). *Let $m = R + 1$ and $\beta = \pi/m$. For $0 < a \leq \beta$ there exists $\mu \in \mathcal{P}(K_a)$ with $E_R(\mu) = 0$. At $a = \beta$ that measure is unique: it is the uniform measure on*

$$Z_m = \left\{ \frac{(2j+1)\pi}{m} : 0 \leq j < m \right\} \pmod{2\pi}. \quad (\text{I.54.1})$$

For $a > \beta$, no probability measure on K_a can have all of its first R moments zero.

Proof. The uniform measure on Z_m has its first $m - 1$ moments zero by summing a geometric progression, and its nearest points to zero have angular distance β . This proves existence for $a \leq \beta$.

For uniqueness at β , define

$$P_0(z) = \sum_{j=0}^{m-2} \sin((j+1)\beta) z^j. \quad (\text{I.54.2})$$

The sine vector is a path eigenvector with eigenvalue $\cos \beta$. Therefore $Q_0(\theta) = (\cos \theta - \cos \beta) |P_0(e^{i\theta})|^2$ has constant Fourier coefficient zero. The calculation in Lemma I.51.2, with $\lambda = 0$ and without normalizing the vector, also shows that all its nonconstant Fourier coefficients are positive. The finite recurrence gives the polynomial identity

$$(1 - 2 \cos \beta z + z^2) P_0(z) = \sin \beta (1 + z^m). \quad (\text{I.54.3})$$

On K_β , the zeros of Q_0 are exactly Z_m : the two roots $e^{\pm i\beta}$ of the denominator in (I.54.3) are the endpoints, and the other $m - 2$ roots of $1 + z^m$ are the zeros of P_0 .

If the first R moments vanish, then $\int Q_0 d\mu = 0$. Since $Q_0 \leq 0$ on K_β , the support lies in Z_m . The Vandermonde moment equations uniquely give the uniform weights. Finally, if $a > \beta$ and $\mu \in \mathcal{P}(K_a)$ had these zero moments, it would also be a measure on K_β and hence this uniform measure. But that measure charges the points $\pm\beta$ inside I_a , a contradiction. $\square$

I.54.2 The other endpoint

At $a = 2\pi/m$, define

$$\nu_m = \frac{1}{m-1} \sum_{j=1}^{m-1} \delta_{2\pi j/m}. \quad (\text{I.54.4})$$

Its first $R = m - 1$ moments all equal $-1/R$. The polynomial coefficients constructed above have finite limits at this endpoint, so (I.50.5) remains valid with $d = 1/m$ and $\rho = 1/R$. The measure (I.54.4) attains the gap bound, and its mixture with δ_0 attains the noisy-gap value for $0 \leq \varepsilon \leq 1/R$. For $m = 2$, this says simply that the open arc I_π is the circle with the antipode removed, so the only measure supported on its complement is δ_π .

There is also uniqueness for the endpoint *empty-gap* optimizer. Equality forces all first R moments to be $-1/R$. Its Toeplitz matrix annihilates $\mathbf{1}$, hence

$$\int \left| 1 + e^{i\theta} + \cdots + e^{iR\theta} \right|^2 d\mu = 0.$$

The support is therefore among the m th roots of unity other than 1. The moment equations give exactly (I.54.4).

We use the endpoint constants only by continuous extension or by these explicit measures. The strict-window root-simplicity and stability statements are not automatically extended to endpoints.

I.54.3 An asymptotically exact frequency–amplitude frontier

Write the normalized gap radius as $\delta = a/(2\pi)$. If $\delta \rightarrow 0$ and the integer cutoff satisfies $R\delta \rightarrow \tau \in (1/2, 1)$, then eventually we are in the first window. Directly from (I.50.4),

$$\frac{\rho_R(2\pi\delta)}{\delta} \rightarrow \mathcal{C}(\tau) := \frac{1}{\tau - \pi^{-1} \tan(\pi\tau)}. \quad (\text{I.54.5})$$

Indeed, $\delta R \rightarrow \tau$, $\delta \cot(\pi\delta) \rightarrow 1/\pi$, and $(R+1)\delta \rightarrow \tau$. This limit is optimal because each finite-parameter value comes with an attaining measure.

The denominator in (I.54.5) decreases from $+\infty$ to 1 as τ increases from $1/2$ to 1; its derivative is $-\tan^2(\pi\tau)$. Thus $\mathcal{C}$ increases from zero to one. The exact tradeoff distinguishes two questions that a single big- O lemma conflates: detecting *any* nonzero coefficient requires about half the inverse gap radius, whereas forcing an asymptotic coefficient of size δ requires a cutoff approaching the full inverse radius within this window.

I.55 A linear-frequency replacement for Lemma 5.2

Source 37, delivered Section 7.

I.55.1 A simple bound without a trigonometric denominator

We now use normalized coordinates $x \in \mathbb{R}/\mathbb{Z}$, with $\hat{\mu}(r) = \int e^{-2\pi i r x} d\mu(x)$.

Theorem I.55.1 (Universal small-arc inequality). *Let $0 < \delta \leq 1/2$ and*

$$R = \lfloor \delta^{-1} \rfloor - 1. \quad (\text{I.55.1})$$

Then every probability measure satisfies

$$\mu\{x : \|x\|_{\mathbb{R}/\mathbb{Z}} < \delta\} \geq \delta - (1 - \delta) \max_{1 \leq r \leq R} |\hat{\mu}(r)|. \quad (\text{I.55.2})$$

If the open arc on the left is empty, there is $1 \leq r \leq R$ with

$$\operatorname{Re} \hat{\mu}(r) \leq -\frac{\delta}{1 - \delta}. \quad (\text{I.55.3})$$

Proof. First let $\delta < 1/2$, and write $\delta^{-1} = m + s$ with $m = \lfloor \delta^{-1} \rfloor \geq 2$ and $0 \leq s < 1$. Then $R = m - 1$ and $\pi < m(2\pi\delta) \leq 2\pi$, so the first-window theorem or its upper endpoint applies. As $m\pi\delta = \pi - s\pi\delta$,

$$\rho_R(2\pi\delta) = \frac{1}{m - 1 + \cot(\pi\delta) \tan(s\pi\delta)}.$$

Convexity of $\tan$ on $[0, \pi/2)$ gives $\tan(sx) \leq s \tan x$ for $0 \leq s \leq 1$. Hence

$$\rho_R(2\pi\delta) \geq \frac{1}{m - 1 + s} = \frac{\delta}{1 - \delta}, \quad d_R(2\pi\delta) \geq \delta. \quad (\text{I.55.4})$$

In (I.50.5), replacing d by the smaller number δ weakens the right side, since $d - (1 - d)E$ is increasing in d for $E \geq 0$. This proves (I.55.2); the one-sided statement follows from (I.50.6) and (I.55.4). At $\delta = 1/2$, use the endpoint $R = 1$, $a = \pi$, or the certificate $(1 + \cos \theta)/2$ directly. $\square$

Translation is harmless. For an arc centered at x_0 , replace $\hat{\mu}(r)$ in the one-sided conclusion by $e^{2\pi i r x_0} \hat{\mu}(r)$. The modulus conclusion is unchanged.

I.55.2 Finite cyclic groups, arbitrary weights, and collisions

Corollary I.55.2 (Weighted finite-cyclic missing interval). *Let N, M be integers with $N \geq 2$, $1 \leq M \leq N/2$, and let $w : \mathbb{Z}/N\mathbb{Z} \rightarrow [0, \infty)$ have positive total mass $W = \sum_x w(x)$. Suppose $w(x) = 0$ whenever the least absolute residue of x is strictly less than M . Then there is an integer r satisfying*

$$1 \leq r \leq \left\lfloor \frac{N}{M} \right\rfloor - 1 < N, \quad -\operatorname{Re} \widehat{w}(r) \geq \frac{M}{N-M} W, \quad (\text{I.55.5})$$

where $\widehat{w}(r) = \sum_{x \in \mathbb{Z}/N\mathbb{Z}} w(x) e^{-2\pi i r x / N}$. The residue class of r is nonzero and has least absolute residue at most the displayed upper bound.

Proof. Apply Theorem I.55.1 to the probability measure that assigns mass $w(x)/W$ to x/N in $\mathbb{R}/\mathbb{Z}$, with $\delta = M/N$. Since $M \geq 1$, the cutoff is at most $N - 1$, so no zero residue is selected. $\square$

[write] **6 October 2026; sources 34–39.** *It implies the formal conclusion* (proved at the write). Under the hypotheses of lemma_5_2 ($1 < N$, $0 < M$, $2M \leq N$; the parity of M is not needed), $A \cap [-M, M) = \emptyset$ implies that $w = 1_A$ vanishes where the least absolute residue is below M . The corollary gives $1 \leq r \leq \lfloor N/M \rfloor - 1 < N$, so $r \neq 0$ in $\mathbb{Z}/N\mathbb{Z}$ and $|r|_N \leq r \leq N/M - 1 \leq (N/M)^2$; and $|\widehat{1_A}(r)| \geq -\operatorname{Re} \widehat{1_A}(r) \geq |A|M/(N-M) \geq |A|M/(4N)$, since $N-M \leq 4N$. The sign convention of the exponential does not affect the modulus. So the conclusion of lemma_5_2 follows, with frequency linear in N/M and coefficient larger by the factor $4N/(N-M) \in (4, 8]$. The formal statement is already exact; its status does not change. The intake checked the corollary exhaustively on $\mathbb{Z}/N\mathbb{Z}$ for $N \leq 14$, every $M \leq N/2$ and every nonempty A avoiding the open gap: 21,781 cases, no failure, 171 equalities. Formalization candidate (the source’s lower-bound track), claimed by nobody: the explicit vector w , its path recurrence and the coefficient identity of Lemma I.51.2.

For an indicator function $w = 1_A$, this gives

$$|\widehat{1_A}(r)| \geq \frac{|A|M}{N-M}. \quad (\text{I.55.6})$$

The original exclusion $A \cap [-M, M) = \emptyset$ implies the open-gap exclusion used here. Thus the corollary applies under the hypotheses of the repository’s Lemma 5.2 and additionally removes the parity assumption on M . If A is empty, the coefficient lower bound is zero and a nonzero frequency exists because $N > 1$; the positive-mass proof is unnecessary.

Quantity	Proof-supported Lemma 5.2	Present corollary
Frequency bound	$(N/M)^2$	$\lfloor N/M \rfloor - 1$
Coefficient bound	$ A M/(4N)$	$ A M/(N-M)$
Coefficient direction	Absolute value	A negative real part
Input	A subset	Any nonnegative weights
Gap assumption	Half-open interval; M even	Open interval; no parity condition

The gain in the coefficient relative to the proof-supported version is $4N/(N-M)$, approaching 4 for a small gap. The frequency range is linear, rather than quadratic, in N/M . The weighted formulation includes empirical measures of polynomial values, with all repetitions counted. Again, the repository already repairs the multiplicity issue; the new part is the bound, not the observation that repeated values must be counted.

I.55.3 Sharpness and endpoint discipline

For $\delta = 1/m$, the measure uniformly supported at j/m , $1 \leq j < m$, has no mass in the *open* arc of radius δ and has $\hat{\mu}(r) = -1/(m-1)$ for $1 \leq r \leq m-1$. Thus the coefficient $\delta/(1-\delta)$ in (I.55.3) is exactly sharp for these radii. When $m \mid N$, this is an indicator example on $\mathbb{Z}/N\mathbb{Z}$ after normalization.

This extremizer charges both endpoints. It therefore need not satisfy the original *half-open* exclusion, which deletes one of them. The exact sharpness assertion is for the open-gap theorem, not for every stronger finite-grid hypothesis. This distinction matters when comparing constants at a fixed small modulus.

The linear order of the cutoff cannot be improved to $o(\delta^{-1})$ in a uniform missing-arc lemma with a positive conclusion. The critical grid in Theorem I.54.1 has a gap radius $1/(2m)$ in normalized coordinates and its first $m-1$ moments vanish. Taking $m \rightarrow \infty$ proves the obstruction. The sharper coefficient-versus-cutoff information is (I.54.5); it is not captured by this order-of-magnitude example.

I.56 Stability and near extremizers

Source 37, delivered Section 8.

I.56.1 A dimension-independent geometric estimate at the threshold

The critical grid has a useful stability property even though the optimal empty-gap coefficient is zero there.

Theorem I.56.1 (Critical geometric stability). *Let $m = R + 1$, $\beta = \pi/m$, and let Z_m be the odd grid in (I.54.1). If $\mu \in \mathcal{P}(K_\beta)$ and $E_R(\mu) \leq \varepsilon$, then*

$$\int \text{dist}(\theta, Z_m)^2 d\mu(\theta) \leq \frac{\pi^2}{m^2} \cot^2\left(\frac{\pi}{2m}\right) \varepsilon \leq 4\varepsilon. \quad (\text{I.56.1})$$

Distance is angular geodesic distance on $\mathbb{T}_{2\pi}$.

Proof. Use P_0, Q_0 from the proof of Theorem I.54.1, and set $D = -Q_0 \geq 0$ on K_β . Since the constant Fourier coefficient of Q_0 is zero and its other coefficients are positive,

$$\int D d\mu \leq Q_0(0)\varepsilon.$$

The sine sum equals $\cot(\beta/2)$, so

$$Q_0(0) = (1 - \cos \beta) \cot^2(\beta/2) = 1 + \cos \beta.$$

Away from the endpoints, (I.54.3) gives

$$D(\theta) = \frac{\sin^2 \beta \cos^2(m\theta/2)}{\cos \beta - \cos \theta} \geq (1 - \cos \beta) \cos^2(m\theta/2).$$

The inequality extends to the endpoints by continuity. Consequently

$$\int \cos^2(m\theta/2) d\mu \leq \cot^2(\beta/2)\varepsilon.$$

If $u = \text{dist}(\theta, Z_m)$, then $0 \leq u \leq \pi/m$ and $\cos^2(m\theta/2) = \sin^2(mu/2) \geq m^2 u^2 / \pi^2$. This proves the first inequality in (I.56.1). The second follows from $x \cot x \leq 1$ for $0 < x \leq \pi/4$. $\square$

The conclusion concerns concentration near the grid, not closeness to the *uniformly weighted* grid in a dimension-independent transport metric. Redistributing the masses among the grid points is a separate issue. At $\varepsilon = 0$, the moment equations do force uniform weights, as already proved.

I.56.2 Quantitative rigidity throughout the strict window

For fixed strict-window parameters, the full extremizer is stable. We give an effective finite-dimensional bound rather than hide the dependence on R and a in an unspecified compactness argument.

Let $Z = \{\theta_1, \dots, \theta_n\}$ be the contact set in (I.53.5), where $n = R + 2$. Let

$$V = (e^{-ik\theta_j})_{-R \leq k \leq R, 1 \leq j \leq n}, \quad \sigma = \sigma_{\min}(V) > 0. \quad (\text{I.56.2})$$

With F from (I.51.5), define

$$h(\theta) = \mathbf{1}_{I_a}(\theta) - F(\theta), \quad \kappa = \inf_{\theta \notin Z} \frac{h(\theta)}{\text{dist}(\theta, Z)^2}. \quad (\text{I.56.3})$$

Lemma I.56.2. *For fixed R, a in the strict window, $\kappa > 0$.*

Proof. The zero set of h is precisely Z . Near zero, $h = 1 - F$ has positive quadratic term $\sum_{k=1}^R k^2 f_k \theta^2$. At every interior point of K_a in Z , the corresponding root of P is simple and $c - \cos \theta > 0$, so $h = -F$ has a positive quadratic zero. At either endpoint, P is nonzero and $\sin a > 0$, so h grows linearly on the outside; on the inside its limit is 1. These observations give a positive lower bound for the ratio in neighborhoods of all contact points. Away from those neighborhoods the function is continuous and positive on a compact set. This proves the claim despite the one-sided jumps at the endpoints. $\square$

Theorem I.56.3 (Stability of the exact mass tradeoff). *Fix $0 \leq \varepsilon \leq \rho_R(a)$ and let μ_ε be the exact extremizer in (I.50.10). Suppose $E_R(\mu) \leq \varepsilon$ and set*

$$\Delta = \mu(I_a) - \frac{\rho_R(a) - \varepsilon}{1 + \rho_R(a)} \geq 0, \quad t = \sqrt{\Delta/\kappa}.$$

Then

$$W_1(\mu, \mu_\varepsilon) \leq t + \frac{\pi\sqrt{R+2}}{2\sigma} \left[2 \sum_{k=1}^R \left(kt + \sqrt{\frac{\varepsilon\Delta}{f_k}} \right)^2 \right]^{1/2}. \quad (\text{I.56.4})$$

In particular the distance is $O_{R,a,\varepsilon}(\sqrt{\Delta})$.

Proof. By the exact slack decomposition (I.53.4), $\int h d\mu \leq \Delta$, whence $\int \text{dist}(\theta, Z)^2 d\mu \leq \Delta/\kappa$. Push each point to a nearest point of Z , with a fixed rule for ties, and call the resulting measure ν . Cauchy–Schwarz gives $W_1(\mu, \nu) \leq t$.

For $z = \hat{\mu}(k)$, the disk constraint $|z| \leq \varepsilon$ gives

$$|z + \varepsilon|^2 \leq 2\varepsilon(\text{Re } z + \varepsilon) \leq \frac{\varepsilon\Delta}{f_k},$$

where the last inequality follows again from (I.53.4). The function $e^{-ik\theta}$ is k -Lipschitz in angular distance, hence

$$|\hat{\nu}(k) - \hat{\mu}_\varepsilon(k)| \leq kt + \sqrt{\varepsilon\Delta/f_k}.$$

Both measures have total mass one, and negative moments are conjugates of positive ones. If u, u^* are their weight vectors on Z , then

$$\|V(u - u^*)\|_2^2 \leq 2 \sum_{k=1}^R (kt + \sqrt{\varepsilon\Delta/f_k})^2.$$

Use $\|u - u^*\|_2 \leq \sigma^{-1} \|V(u - u^*)\|_2$ and $\text{TV}(\nu, \mu_\varepsilon) \leq \sqrt{R+2} \|u - u^*\|_2/2$. Since the diameter of the circle is π , $W_1(\nu, \mu_\varepsilon) \leq \pi \text{TV}(\nu, \mu_\varepsilon)$. The triangle inequality completes the proof. $\square$

The coefficients f_k , the nodes, and the Vandermonde matrix are explicit. The constant κ is a one-variable positive minimum with explicitly controlled local behavior. The estimate is therefore suitable for certified parameter-specific bounds. We do not assert uniform control of κ or σ as the cutoff grows or as the endpoints of the window are approached.

I.57 A stronger monomial-recurrence deduction

Source 37, delivered Section 9.

The sharp spectral lemma can be used without its full extremal formula. The simple conclusion (I.55.3) is enough. A second improvement is to choose the Dirichlet denominator scale according to the size of the detected Weyl sum, rather than fixing it at the sample length.

I.57.1 The external exponential-sum input

Write $e(x) = e^{2\pi ix}$ and

$$S_k(\alpha; T) = \sum_{n=1}^T e(\alpha n^k).$$

For fixed integers $k \geq 2$, $T \geq 1$, and real parameters $D \geq 1$, $B \geq 1$, consider the hypothesis

$$\frac{|S_k(\beta; T)|}{T} \leq B \left(q^{-1} + T^{-1} + qT^{-k} \right)^{1/D} \quad (\text{I.57.1})$$

for all $\beta \in \mathbb{R}$ and all coprime integers u, q with $q > 0$ and $|\beta - u/q| \leq q^{-2}$. It is important that the hypothesis permits any positive denominator; no restriction $q \leq T$ is imposed.

Theorem I.57.1 (A recurrence conversion with optimized denominator). *Assume (I.57.1). If $0 < \delta \leq 1/2$ and*

$$T\delta^{D+1} \geq 16B^D, \quad (\text{I.57.2})$$

then every $\alpha \in \mathbb{R}$ admits an integer $1 \leq p \leq T$ such that

$$\left\| \alpha p^k \right\|_{\mathbb{R}/\mathbb{Z}} < \delta. \quad (\text{I.57.3})$$

Proof. Suppose the conclusion fails. The probability measure

$$\mu = \frac{1}{T} \sum_{n=1}^T \delta_{\alpha n^k \bmod 1}$$

has no mass in the open arc of radius δ . This is an empirical measure, not a set of distinct values; repetitions are counted. Theorem I.55.1 provides

$$1 \leq r \leq \lfloor \delta^{-1} \rfloor - 1 < \delta^{-1}, \quad \frac{|S_k(r\alpha; T)|}{T} \geq \frac{\delta}{1 - \delta} \geq \delta. \quad (\text{I.57.4})$$

Put

$$s = (\delta/B)^D, \quad Q = \left\lfloor \frac{T^k s}{4} \right\rfloor. \quad (\text{I.57.5})$$

The assumption says $T\delta s \geq 16$, and therefore $Ts \geq 32$. Since $k \geq 2$, the expression inside the floor is at least 2, so

$$\frac{T^k s}{8} \leq Q \leq \frac{T^k s}{4}, \quad Q \geq 1. \quad (\text{I.57.6})$$

Dirichlet approximation gives coprime u, q , $1 \leq q \leq Q$, with

$$|r\alpha - u/q| \leq \frac{1}{qQ} \leq \frac{1}{q^2}. \quad (\text{I.57.7})$$

For clarity, this version follows by placing the $Q + 1$ fractional parts of $0, \beta, \dots, Q\beta$ in Q equal intervals and subtracting two in the same interval; reducing the resulting fraction only improves the displayed bound with its reduced denominator.

Apply (I.57.1) to $r\alpha$. By (I.57.4),

$$s \leq q^{-1} + T^{-1} + qT^{-k}.$$

We have $qT^{-k} \leq s/4$, and $T^{-1} \leq \delta s/16 \leq s/16$. Thus $q^{-1} > s/2$, and in particular $q \leq 2/s$. Set $p = rq$. It is a positive integer, and

$$p < \frac{2}{\delta s} = \frac{2B^D}{\delta^{D+1}} \leq \frac{T}{8}. \quad (\text{I.57.8})$$

Multiplying (I.57.7) by p^k/r shows that αp^k lies within p^{k-1}/Q of the integer $ur^{k-1}q^{k-1}$. Hence

$$\left\| \alpha p^k \right\|_{\mathbb{R}/\mathbb{Z}} \leq \frac{p^{k-1}}{Q} \leq \frac{8}{Ts} \leq \frac{\delta}{2} < \delta.$$

This contradicts the assumed absence of recurrence and proves the theorem. $\square$

The use of $Q \asymp T^k s$ is legitimate because the rational approximation still meets the q^{-2} condition in the Weyl input. It is precisely the smallness of the term qT^{-k} that allows this larger search scale.

I.57.2 Power-law consequences

Corollary I.57.2 (General power-law Weyl input). *Suppose (I.57.1) holds for all integers $T \geq T_0$ with $B = CT^\eta$, where $C \geq 1$ and $0 \leq \eta < 1/D$. Define*

$$A = (16C^D)^{1/(D+1)}, \quad b = \frac{1 - D\eta}{D + 1}. \quad (\text{I.57.9})$$

Then for every $T \geq T_0$ and every $\alpha \in \mathbb{R}$,

$$\min_{1 \leq p \leq T} \|\alpha p^k\|_{\mathbb{R}/\mathbb{Z}} \leq AT^{-b}. \quad (\text{I.57.10})$$

Proof. Set $\delta = AT^{-b}$. If $\delta \geq 1/2$, the bound is trivial. Otherwise

$$T\delta^{D+1} = 16C^D T^{D\eta} = 16B^D,$$

and Theorem I.57.1 applies. □

For the classical all- η Weyl inequality, $D = 2^{k-1}$. Choosing $\eta > 0$ arbitrarily small gives the transparent local consequence

$$\min_{1 \leq p \leq T} \|\alpha p^k\|_{\mathbb{R}/\mathbb{Z}} \ll_{k,\epsilon} T^{-1/(2^{k-1}+1)+\epsilon} \quad (\epsilon > 0). \quad (\text{I.57.11})$$

The point here is the deduction and its explicit dependence on the Weyl input. No claim is made that (I.57.11) is a new record.

I.57.3 Using the explicit input recorded in ProveIt

The explicit version of the Weyl input in the inspected Section 5 source and its companion lemma_5_3_holds has

$$D = 2^{k-1}, \quad C = 1000, \quad \eta = \frac{1}{k2^{k+1}} = \frac{1}{4kD}, \quad T_0 = 2^{2^{40k^3}}. \quad (\text{I.57.12})$$

We use this statement as an external input; its proof is not reproduced here, and the Lean project was not built during preparation of this manuscript. The displayed companion was inspected in `Proofs05Weyl.lean` at the pinned revision [167].

Corollary I.57.3 (Explicit refinement of the monomial recurrence step). *For $k \geq 2$, let*

$$A_k = (16 \cdot 1000^{2^{k-1}})^{1/(2^{k-1}+1)}, \quad b_k = \frac{1 - 1/(4k)}{2^{k-1} + 1}. \quad (\text{I.57.13})$$

Using the input (I.57.12), for every integer $T \geq 2^{2^{40k^3}}$ and every $\alpha \in \mathbb{R}$ there is $1 \leq p \leq T$ with

$$\|\alpha p^k\|_{\mathbb{R}/\mathbb{Z}} \leq A_k T^{-b_k}. \quad (\text{I.57.14})$$

For an integer $N \geq 2$ and $a \in \mathbb{Z}/N\mathbb{Z}$, the corresponding least-residue bound is $|p^k a|_N \leq A_k T^{-b_k} N$.

Proof. Substitute (I.57.12) into Corollary I.57.2. For the finite-cyclic assertion take α to be any integer representative of a divided by N . □

[write] **6 October 2026; sources 34–39.** *It implies the formal bound of Lemma 5.5* (proved at the write). The input (I.57.12) is the second conjunct of `lemma_5_3` exactly: the constant 1000, the exponent $1/(k2^{k+1})$, any $q > 0$ coprime to a with $|\alpha - a/q| \leq q^{-2}$, `weylSum` summing over $1 \leq s \leq t$, and $t \geq \text{weylThreshold } k = 2^{2^{40k^3}}$. For such t , $A_k t^{-b_k} \leq t^{-a_k}$ because $A_k < 16 \cdot 1000$ while $t^{b_k - a_k} \geq 2^{2^{40k^3}(b_k - a_k)}$ is far larger; so the corollary gives the conclusion of `lemma_5_5` (with $p \leq t$, $|p^k a|_N \leq t^{-a_k} N$). The comparison is with one local step at its chosen Weyl parameters, as the source says; source 32’s explicit quadratic recurrence $p \leq 2^{10} R^5$ (Lemma I.45.3, every scale, small constants) and the Lau- and Maynard-based recurrences of sources 13, 16, 18 and 19 are other regimes. Formalization candidate, claimed by nobody: Dirichlet at scale $\lfloor T^k s/4 \rfloor$ with Theorem I.55.1.

The recurrence exponent in the selected Lemma 5.5 is $a_k = 1/(k2^{k+1}) = 1/(4kD)$. The new exponent is larger by the factor

$$\frac{b_k}{a_k} = \frac{(4k-1)2^{k-1}}{2^{k-1} + 1}. \quad (\text{I.57.15})$$

The improvement comes with the explicitly displayed multiplicative constant A_k ; it is not a constant-free comparison for every small T .

Degree k	Selected exponent a_k	Exponent b_k here	b_k/a_k
2	1/16	7/24	14/3
3	1/48	11/60	44/5
4	1/128	5/48	40/3
5	1/320	19/340	304/17
6	1/768	23/792	736/33

This is a comparison against a particular local step, with exactly its chosen Weyl parameters. It is not a comparison against the strongest available bounds for fractional parts of monomials or general polynomials.

I.57.4 The square-root-length interface

The degree induction in a polynomial partition needs a short step as well as a small leading coefficient. Integer rounding can be kept explicit.

Corollary I.57.4 (Short-step recurrence). *Under the hypotheses of Corollary I.57.2, let $U \geq 4$ be an integer with $\lfloor \sqrt{U} \rfloor \geq T_0$. There is a positive integer p satisfying*

$$p^2 \leq U, \quad \left\| \alpha p^k \right\|_{\mathbb{R}/\mathbb{Z}} \leq 2^b A U^{-b/2}. \quad (\text{I.57.16})$$

Proof. Apply (I.57.10) with $T = \lfloor \sqrt{U} \rfloor$. Then $p \leq T \leq \sqrt{U}$, while $T \geq \sqrt{U}/2$ because $U \geq 4$. $\square$

[write] **6 October 2026; sources 34–39.** The hypothesis $\lfloor \sqrt{U} \rfloor \geq T_0$ is the rounding-safe form of the squared threshold of the formal `lemma_5_5_square_root_auxiliary` (`Section05.lean`, line 192: `weylThreshold` $k^2 \leq t, p^2 \leq t$), as the source says in Section F.42.

For (I.57.12), this gives exponent $b_k/2$, rather than silently claiming exponent b_k at a square-root sample length. Since $b_k/2 > a_k$ for every $k \geq 2$, it also recovers the original auxiliary target U^{-a_k} whenever, in addition,

$$U \geq (2^{b_k} A_k)^{1/(b_k/2 - a_k)}. \quad (\text{I.57.17})$$

The explicit square-root threshold and the constant-absorption threshold must both be propagated in any replacement proof of a partition lemma.

I.57.5 Where the improvement stops in this manuscript

The conclusions above concern *monomial* recurrence. A bound for αp^k does not by itself make the lower coefficients of an arbitrary polynomial small. The partition induction must still control those terms, its target lengths, and the losses from refinement. The existing source also records corrections to exact-cell formulations. None of those issues is bypassed by an improved leading-coefficient estimate.

Likewise, a better local recurrence exponent is not an improved global Szemerédi bound until every later parameter has been recomputed. The present work establishes the spectral and recurrence substitutions, gives the precise short-step interface, and leaves that global propagation as a separate research task.

I.58 The real interval factorization used above

Source 37, delivered Appendix A.

We include the elementary factorization argument behind Lemma I.52.1. This avoids treating supported moment existence as a black box.

Lemma I.58.1 (Real spectral factorization). *If a real even trigonometric polynomial $q(\phi)$ of degree at most n is nonnegative for all ϕ , then*

$$q(\phi) = |D(e^{i\phi})|^2$$

for a polynomial D of degree at most n with real coefficients.

Proof. The zero polynomial is immediate. Otherwise, let $d \leq n$ be the actual trigonometric degree and write $q(\phi) = \sum_{j=-d}^d c_j z^j$ on $z = e^{i\phi}$. Since q is real and even, $c_j = c_{-j} \in \mathbb{R}$. Thus $H(z) = z^d q(\phi)$ is a real reciprocal polynomial of degree $2d$, with nonzero constant term. Its zeros off the unit circle occur with their reciprocals and conjugates. A zero on the unit circle has even multiplicity: otherwise the real analytic nonnegative function $q(\phi)$ would change sign at that zero. The local change of variable $z = e^{i\phi}$ has nonzero derivative, so the multiplicities agree.

Select from each reciprocal pair off the unit circle the zeros inside the circle, respecting conjugate pairs, and select half of each unit-circle multiplicity, again respecting conjugation. The resulting monic factor D_0 has degree d and real coefficients. Its reciprocal polynomial $D_0^*(z) = z^d D_0(1/z)$ supplies the complementary zeros. Hence $H(z) = c D_0(z) D_0^*(z)$ for a real constant c . On the unit circle, $z^{-d} D_0(z) D_0^*(z) = |D_0(z)|^2$. At a point where q is positive this shows $c > 0$. Taking $D = \sqrt{c} D_0$ proves the assertion. $\square$

Lemma I.58.2 (Interval factorization). *If p is a real polynomial of degree at most $2R$ and is nonnegative on $[-b, b]$, where $b > 0$, then*

$$p(t) = U(t)^2 + (b^2 - t^2)V(t)^2 \tag{I.58.1}$$

for real polynomials with $\deg U \leq R$ and $\deg V \leq R - 1$.

Proof. Apply Lemma I.58.1 to $q(\phi) = p(b \cos \phi)$, which has degree at most $2R$. Let $D(z) = \sum_{j=0}^{2R} d_j z^j$ be the resulting real factor, with zeros added to the coefficient list if its degree is smaller. The Laurent polynomial $e^{-iR\phi} D(e^{i\phi})$ has frequencies between $-R$ and R . Its real part is a real polynomial $A(\cos \phi)$ of degree at most R , because each $\cos(j\phi)$ is such a polynomial. Its imaginary part is $\sin \phi B(\cos \phi)$ with $\deg B \leq R-1$, because $\sin(j\phi)/\sin \phi$ is a polynomial of degree $|j|-1$ for nonzero integer j . Therefore

$$p(b \cos \phi) = A(\cos \phi)^2 + (1 - \cos^2 \phi) B(\cos \phi)^2.$$

Set $U(t) = A(t/b)$ and $V(t) = B(t/b)/b$. This proves (I.58.1) on the whole interval and hence as a polynomial identity on $\mathbb{R}$. $\square$

I.59 Low-dimensional examples

Source 37, delivered Appendix B.

I.59.1 One retained frequency

For $R = 1$, the first window is $\pi/2 < a < \pi$. Let $c = \cos a < 0$. Then

$$P(z) = 1, \quad \rho = -c, \quad d = \frac{-c}{1-c}, \quad F(\theta) = \frac{\cos \theta - c}{1-c}.$$

The unique empty-gap extremizer is $(\delta_{-a} + \delta_a)/2$. For $0 \leq \varepsilon \leq -c$, the exact minimum gap mass is $(-c - \varepsilon)/(1 - c)$, attained by adding that mass at zero and rescaling the two endpoint atoms. This example checks the sign, normalization, and open-endpoint conventions directly.

I.59.2 Two retained frequencies

For $R = 2$, the first window is $\pi/3 < a < 2\pi/3$, so $-1/2 < c = \cos a < 1/2$. Symmetry gives

$$\begin{aligned} P(z) &= \frac{1+z}{2}, & F(\theta) &= \frac{(\cos \theta - c)(1 + \cos \theta)}{2(1-c)}, \\ \rho &= \frac{1-2c}{3-2c}, & d &= \frac{1-2c}{4(1-c)}. \end{aligned}$$

The support is $\{-a, a, \pi\}$. The mass at each endpoint is

$$u = \frac{1}{(3-2c)(1+c)},$$

and the mass at π is

$$v = 1 - 2u = \frac{(1-c)(1+2c)}{(3-2c)(1+c)}.$$

All masses are positive in the strict window. The two required identities are

$$2uc - v = -\rho, \quad 2u(2c^2 - 1) + v = -\rho.$$

They are included as exact symbolic checks in `verify.py`. At $a = \pi/2$, the three masses are equal, $\rho = 1/3$, and $d = 1/4$. Thus zero first and second moments force at least $1/4$ of the mass into the open arc $(-\pi/2, \pi/2)$, with equality at the uniform measure on the fourth roots of unity. This is a concrete instance of the general noisy-gap extremizer at $\varepsilon = 0$.

I.59.3 Practical recovery of an extremizer

For a strict-window pair (R, a) , the following steps implement the exact mathematical construction numerically.

1. Compute S, d, ρ and the coefficients w_j from (I.50.3)–(I.50.11).
2. Find the $R - 1$ roots of $P(z)$ and combine their arguments with $-a, a$. The proof guarantees distinct unit-circle roots on the allowed arc.
3. Solve (I.50.13), or its overdetermined version using frequencies $-R, \dots, R$. Check positivity and moment residuals.
4. For a chosen $\varepsilon \leq \rho$, add the atom at zero with mass $(\rho - \varepsilon)/(1 + \rho)$ and rescale the other masses.

Near a window endpoint, direct trigonometric evaluation can lose relative accuracy. A production implementation should use stable limiting formulas or higher precision, and certified interval arithmetic for a machine-checkable numerical artifact. The included double-precision script is a regression suite, not that certification layer.

I.60 Source 35’s Part 3: polynomial divisibility and prescribed lengths [write]

Added in the seventh write, 6 October 2026.

The next section is source 35’s Section 7, its delivered Part 3, on partitions of an interval into progressions of prescribed lengths on which several polynomial phases are nearly constant (Gowers’s Lemma 5.9). *Logically it follows source 23’s chapter* (Sections I.35–I.39, the linear mixed-radix construction and the numerical-semigroup thresholds) *and source 19’s polynomial obstruction* (Theorem I.34.7); it is printed here because no number printed earlier may move. Source 35’s introduction is Section F.40; its Parts 1 and 2 are in Parts II and III.

What is proved. Assigning a different prime denominator to each degree, explicit phases with degrees $d_1, \dots, d_\nu$ force every flat progression of length at least L to have step divisible by an explicit $D \asymp L^W \prod_i \varepsilon_i^{-d_i}$, $W = \sum_i d_i(d_i + 1)/2$, whatever its starting point (Theorem I.61.2); the exact full-partition threshold for a set of lengths with conductor $C(\mathcal{S})$ is then $DC(\mathcal{S})$ (Theorem I.61.4); for the lengths L or $L + 1$, every target, the necessary frontier is $(W + 2)a + \sum_i d_i b_i \leq 1$ (Corollary I.61.6), that is $1/(7\nu + 2)$ for quadratics at $b = 2a$ and source 23’s sharp $1/(3\nu + 2)$ in degree one. This is the higher-degree analogue of source 23’s construction (its Question 6, item I.15); the matching sufficient theorem in degree at least two is open. The flexible-length obstruction of source 19 is unchanged (note after the corollary).

Relation to the formal statements (checked at HEAD 89c1cbc38). `lemma_5_9` (Section05.lean, line 261) has the exact companion `lemma_5_9_holds` (Proofs05FullPartition.lean, line 98). The results here are obstructions, lower bounds on what any prescribed-length theorem can give; nothing is formalized.

I.61 Polynomial divisibility and prescribed progression lengths

Source 35, delivered Section 7.

The loss caused by prescribing progression lengths is already understood for linear phases in the preceding report [165]. Its construction forces every sufficiently long flat progression to have a common divisor in its step, after which a numerical-semigroup obstruction prevents an exact partition. The higher-degree version was left as a research question. We resolve that construction problem here. Distinct prime denominators at successive polynomial degrees force successive divisibility conditions; their product has the weighted-degree scale suggested by the earlier volume argument [164].

The resulting obstruction has two consequences. It excludes long flat progressions at *every* initial point, whereas the volume argument only needed to exclude a cell through one endpoint. For prescribed lengths, the period is multiplied by the conductor of the allowed lengths. This strengthens the necessary exponent region for that interface of Gowers' Lemma 5.9 [1]. The construction does not provide a matching upper bound for general polynomial phases of degree at least two.

I.61.1 A small-arc estimate with an explicit constant

Write $\mathbb{T} = \mathbb{R}/\mathbb{Z}$. The arc diameter of a finite subset of $\mathbb{T}$ is the length of its shortest containing closed circular arc. A proper integer progression has the form

$$P = \{a, a + h, \dots, a + (\ell - 1)h\}, \quad h \in \mathbb{Z}_{>0}.$$

All progressions in this section are understood in this sense, including when the phases themselves take values in a cyclic group.

Lemma I.61.1 (Leading coefficient of a phase in a short arc). *Let $j \geq 1$, $L \geq j + 1$, and $0 < \varepsilon < 2^{-j}$. Suppose a real polynomial $P(t)$ of degree at most j , with coefficient α_j of t^j , satisfies*

$$\text{diam}_{\text{arc}}\{P(0), \dots, P(L - 1)\} \leq \varepsilon \quad \text{in } \mathbb{T}.$$

With $H_j = \lfloor (L - 1)/j \rfloor$, one has

$$\|j! \alpha_j\|_{\mathbb{T}} \leq \frac{2^{j-1} \varepsilon}{H_j^j}. \quad (\text{I.61.1})$$

Proof. Choose real lifts u_n of $P(n)$ in one interval of length ε . Every available $(j + 1)$ st forward difference of this sequence is an integer, because the corresponding difference of P is zero. The positive and negative binomial weights each sum to 2^j . Consequently

$$|\Delta^{j+1} u_n| \leq 2^j \varepsilon < 1,$$

so all these differences vanish. Newton interpolation therefore gives a real polynomial Q of degree at most j with $Q(n) = u_n$ throughout the given interval. If $L = j + 1$, ordinary interpolation gives the same conclusion directly, without a difference condition.

Let β_j be the coefficient of t^j in Q . Since $u_n - P(n)$ is integral,

$$j! \beta_j = \Delta_1^j Q(0) \equiv \Delta_1^j P(0) = j! \alpha_j \pmod{1}.$$

The points $0, H_j, \dots, jH_j$ all lie in $[0, L - 1]$. Taking a difference with this larger step gives

$$|H_j^j j! \beta_j| = |\Delta_{H_j}^j Q(0)| \leq 2^{j-1} \varepsilon.$$

The last inequality again uses the positive and negative binomial weights, whose sums are now 2^{j-1} . Dividing by H_j^j and using the congruence proves (I.61.1). $\square$

This is a concrete polynomial smoothness estimate. Such estimates are part of the broader quantitative theory of polynomial orbits [148]; the elementary proof above is all that is needed here. The strict small-arc hypothesis is what permits the modular differences to be lifted without an integer ambiguity.

I.61.2 One prime denominator at each degree

To avoid confusing the number of phases with the field size in the preceding sections, write ν for the number of phases. Fix positive degrees $d_1, \dots, d_\nu$, and put

$$d_* = \max_i d_i, \quad J = \sum_{i=1}^{\nu} d_i, \quad W = \sum_{i=1}^{\nu} \frac{d_i(d_i + 1)}{2}. \quad (\text{I.61.2})$$

Thus J counts the nonconstant coefficients and W is their total degree weight.

Theorem I.61.2 (Exact divisibility of all sufficiently long flat cells). *Let $L \geq d_* + 1$ and $0 < \varepsilon_i < 2^{-d_i}$. For every pair (i, j) with $1 \leq i \leq \nu$ and $1 \leq j \leq d_i$, choose a prime $p_{ij} > j$. Assume all these primes are distinct and*

$$p_{ij} 2^{j-1} \varepsilon_i < \left\lfloor \frac{L-1}{j} \right\rfloor^j. \quad (\text{I.61.3})$$

Define

$$\Phi_i(x) = \sum_{j=1}^{d_i} \frac{x^j}{p_{ij}} \pmod{1}, \quad D = \prod_{i=1}^{\nu} \prod_{j=1}^{d_i} p_{ij}. \quad (\text{I.61.4})$$

For every integer progression P of length at least L and positive step h , the following conditions are equivalent:

1. $\text{diam}_{\text{arc}} \Phi_i(P) \leq \varepsilon_i$ for every i ;
2. D divides h .

In the second case every Φ_i is constant on P .

Proof. Assume the diameter bounds and retain the first L points of P . Fix i . We prove $p_{ij} \mid h$ by descending induction on j . Suppose this has been proved for every larger degree t . Then

$$\frac{(a + nh)^t - a^t}{p_{it}} \in \mathbb{Z} \quad (n \in \mathbb{Z}),$$

because $p_{it} \mid h$. Thus each term of degree greater than j is constant modulo one along the progression. Subtracting these constants leaves a degree- j polynomial in n whose leading coefficient is h^j/p_{ij} and whose arc diameter is still at most ε_i .

Lemma I.61.1, applicable since $\varepsilon_i < 2^{-d_i} \leq 2^{-j}$, gives

$$\left\| \frac{j! h^j}{p_{ij}} \right\|_{\mathbb{T}} \leq \frac{2^{j-1} \varepsilon_i}{\lfloor (L-1)/j \rfloor^j} < \frac{1}{p_{ij}}.$$

If $p_{ij} \nmid h$, then $p_{ij} > j$ makes $j! h^j$ nonzero modulo p_{ij} . Its distance, divided by p_{ij} , from the nearest integer is at least $1/p_{ij}$, a contradiction. This proves $p_{ij} \mid h$ and completes the descending induction. Repeating it for every phase and using the distinctness of the primes proves $D \mid h$.

Conversely, if $D \mid h$, then for every term in (I.61.4) the difference $(a + nh)^j - a^j$ is divisible by its denominator. Every phase is therefore constant on P . $\square$

The order of the induction matters. A lower-degree coefficient of $\Phi_i(a + nh)$ generally contains contributions from all higher degrees. Once a higher denominator divides h , its *entire* term becomes constant modulo one along the progression, so those contributions disappear before the next degree is examined.

Lemma I.61.3 (Choosing distinct denominators). *Set*

$$X_{ij} = \frac{\lfloor (L-1)/j \rfloor^j}{2^{j-1}\varepsilon_i}. \quad (\text{I.61.5})$$

If $X_{ij}/2^J > \max(d_*, 2)$ for every pair (i, j) , primes satisfying Theorem I.61.2 can be chosen so that

$$2^{-J^2} \prod_{i,j} X_{ij} < D < \prod_{i,j} X_{ij}. \quad (\text{I.61.6})$$

Proof. For a fixed $X = X_{ij}$, the J disjoint intervals

$$(X/2^{u+1}, X/2^u), \quad 0 \leq u < J,$$

each contain a prime by Bertrand's postulate. Its usual integer form suffices: for a real lower endpoint $y > 2$, apply it to $\lfloor y \rfloor$; the resulting prime lies strictly between y and $2y$. All these primes exceed d_* . Choose the pairs (i, j) in any order. At most $J - 1$ primes have previously been used, so one of the J candidates is new. It lies between $X_{ij}/2^J$ and X_{ij} and therefore satisfies (I.61.3). Multiplying the J resulting inequalities proves (I.61.6). $\square$

In particular, for a fixed degree list and sufficiently large L , this choice is available uniformly for all tolerances in the stipulated range. An explicit lower constant is useful. If $L \geq 2d_* + 2$, then $\lfloor (L-1)/j \rfloor \geq L/(2j)$, and

$$D > c_d L^W \prod_{i=1}^{\nu} \varepsilon_i^{-d_i}, \quad c_d = 2^{-J^2 - \sum_i d_i^2} \prod_{i=1}^{\nu} \prod_{j=1}^{d_i} j^{-j}. \quad (\text{I.61.7})$$

Here and below the explicit hypothesis in Lemma I.61.3 is retained whenever a finite bound is used. The corresponding upper bound $D < L^W \prod_i \varepsilon_i^{-d_i}$ follows immediately from (I.61.6). No optimality is claimed for the constants in (I.61.7).

I.61.3 Exact partition thresholds and exceptional sets

Let $\mathcal{S}$ be a nonempty finite set of permitted positive lengths, with $\min \mathcal{S} = L$. Write

$$\langle \mathcal{S} \rangle = \left\{ \sum_{u \in \mathcal{S}} n_u u : n_u \in \mathbb{Z}_{\geq 0} \right\}.$$

If $\gcd(\mathcal{S}) = 1$, its conductor $C(\mathcal{S})$ is the least integer C for which every integer $n \geq C$ belongs to this semigroup. In particular $C - 1 \notin \langle \mathcal{S} \rangle$. We use the term *good* for progressions satisfying all the diameter bounds in Theorem I.61.2.

Theorem I.61.4 (Exact thresholds for the constructed phases). *For the phases in Theorem I.61.2:*

1. *The least length of an integer interval containing a good progression of length L is exactly $D(L-1) + 1$. This is independent of the initial point of the interval.*

2. If $\gcd(\mathcal{S}) = 1$, the least integer R_0 such that every interval of every length $R \geq R_0$ has a complete partition into good progressions of lengths in $\mathcal{S}$ is exactly

$$R_0 = DC(\mathcal{S}). \quad (\text{I.61.8})$$

The statement permits different cells to have different steps.

Proof. Every good progression of length at least L has step at least D . Its integer span is therefore at least $D(L - 1)$. Conversely, a progression of length L and step D is good and can start at the first point of any interval with at least $D(L - 1) + 1$ points.

For the second assertion, every good cell lies in one residue class modulo D . An interval of length $DC - 1$ has $D - 1$ residue classes with C points and one with $C - 1$ points. A partition would express the size of the latter class as a sum of lengths in $\mathcal{S}$, contradicting the definition of the conductor.

If $R \geq DC$, every residue class in an interval of length R has at least C points. Express its size as a sum of permitted lengths and split that residue chain into consecutive blocks of those sizes. Every block has step D and is good. Doing this for all residue classes gives the required complete partition. $\square$

The same argument gives the full packing profile, including sizes below the conductor. Define

$$s_{\mathcal{S}}(n) = \max(\langle \mathcal{S} \rangle \cap [0, n]), \quad g_{\mathcal{S}}(n) = n - s_{\mathcal{S}}(n). \quad (\text{I.61.9})$$

The maximum exists because zero is allowed in the semigroup.

Proposition I.61.5 (Exact uncovered mass). *Write $R = aD + b$ with $0 \leq b < D$. The minimum number of uncovered points in a packing of an interval of length R by disjoint good progressions with lengths in $\mathcal{S}$ is*

$$(D - b)g_{\mathcal{S}}(a) + bg_{\mathcal{S}}(a + 1). \quad (\text{I.61.10})$$

This remains valid without the assumption $\gcd(\mathcal{S}) = 1$.

Proof. There are $D - b$ residue classes of size a and b of size $a + 1$. All cells remain within individual classes. The number of covered points in a class of size n is a semigroup element at most n , and is therefore at most $s_{\mathcal{S}}(n)$. This proves the lower bound on uncovered points. To attain it, use a semigroup representation of $s_{\mathcal{S}}(n)$ and take consecutive blocks of the corresponding lengths in that residue chain. All have step D and are good. $\square$

Thus the exceptional-set issue is also explicit for these examples. If $\gcd(\mathcal{S}) > 1$, arbitrarily large residue-class sizes fail the divisibility condition for membership in the semigroup, and no finite full-partition threshold exists for these phases.

I.61.4 Consequences for universal polynomial partition bounds

Let $\mathcal{T}_{\mathcal{S}, \mathbf{d}}(\varepsilon)$ denote the least threshold after which *every* family of ν real polynomial phases of degrees at most d_i admits the prescribed partition on every integer interval. Allow the value $+\infty$ if there is no such threshold. Theorems I.61.2 and I.61.4 give the unconditional lower bound

$$\mathcal{T}_{\mathcal{S}, \mathbf{d}}(\varepsilon) \geq DC(\mathcal{S}) > c_{\mathbf{d}}C(\mathcal{S})L^W \prod_i \varepsilon_i^{-d_i}. \quad (\text{I.61.11})$$

This combines two distinct constraints: polynomial flatness forces the step divisor D , and exact coverage forces the conductor.

For two consecutive permitted lengths, the elementary Frobenius calculation gives

$$C(\{L, L+1\}) = L(L-1). \quad (\text{I.61.12})$$

Indeed, write $n = aL + b$, $0 \leq b < L$. If $n \geq L(L-1)$, then $a \geq L-1 \geq b$ and $n = (a-b)L + b(L+1)$. On the other hand, a representation $L(L-1) - 1 = uL + v(L+1)$ would force $v \equiv L-1 \pmod{L}$, and hence $v(L+1) \geq (L-1)(L+1) > L(L-1) - 1$.

Consequently (I.61.11) has order at least

$$L^{W+2} \prod_i \varepsilon_i^{-d_i} \quad \text{for } \mathcal{S} = \{L, L+1\}. \quad (\text{I.61.13})$$

For equal degrees $d_i = k$ and equal tolerances $\varepsilon_i = \varepsilon$, this is

$$L^{\nu k(k+1)/2+2} \varepsilon^{-k\nu}.$$

Corollary I.61.6 (Necessary exponents for prescribed lengths). *Fix $c > 0$ and $C_i > 0$. Suppose a universal assertion supplies, for every sufficiently large integer R and every integer target L with $2 \leq L \leq cR^a$, a complete partition into lengths L or $L+1$, for arbitrary phases of degrees at most d_i , with diameters at most $C_i R^{-b_i}$. Assume $a > 0$ and $b_i \geq 0$; if $b_i = 0$, require $C_i < 2^{-d_i}$. Then necessarily*

$$(W+2)a + \sum_i d_i b_i \leq 1. \quad (\text{I.61.14})$$

For equal degrees and tolerances, this becomes

$$\left(\frac{\nu k(k+1)}{2} + 2 \right) a + k\nu b \leq 1. \quad (\text{I.61.15})$$

Proof. Fix a large target L and let $R_L = \lceil (L/c)^{1/a} \rceil$, increased by a fixed amount if needed to enter the asserted large- R range. Put $\varepsilon_i = C_i R_L^{-b_i}$. The universal assertion, applied to this same target for every $R \geq R_L$, gives a partition with tolerance at most ε_i , since $b_i \geq 0$. Hence

$$\mathcal{T}_{\{L, L+1\}, \mathbf{d}}(\varepsilon) \leq R_L.$$

For large L the small-arc conditions hold. Combining this upper bound with (I.61.11) and (I.61.12) gives, up to fixed positive constants,

$$L^{1/a} \geq L^{W+2+\sum_i d_i b_i/a}.$$

Letting L tend to infinity proves (I.61.14). □

[write] **6 October 2026; sources 34–39.** *In this report* (checked at the write). The corollary concerns partitions into the prescribed lengths L or $L+1$ for every target L , the quantifier of Gowers’s Lemma 5.9. Source 19’s Theorem I.34.7 ($(S+1)a + Db \leq 1$; at $b = 2a$ for quadratics $1/(7\nu+1)$) is an obstruction for partitions with flexible lengths at least r^a , which implies one for prescribed lengths; the conductor contributes one more power of L here, giving $(W+2)a + \sum_i d_i b_i \leq 1$, $W = S$ (at $b = 2a$: $a \leq 1/(\nu k(k+5)/2+2)$, which is $1/(7\nu+2)$ for quadratics and source 23’s sharp $1/(3\nu+2)$ for $k=1$, Proposition I.38.1). The flexible-length obstruction of source 19 is unchanged. With Theorem I.61.2 this answers the construction half of source 23’s Question 6 (a higher-degree analogue of its mixed-radix construction; item I.15); the matching necessity, that

is a sufficient theorem in degree at least two, is open (Research question V.133.9). The intake checked the conductor formula for $L \leq 14$ and, by brute force over 1,162, 10,680, 14,220 and 80,968 progressions for the degrees 1, 2, 2, 3 (prime lists (7), (3, 5), (5, 3), (5, 7, 11)), that flat progressions are exactly those with $D \mid h$.

The quantifier over all target lengths is the one relevant to the prescribed-length conclusion of Gowers' Lemma 5.9. Its role in the proof is essential: a fixed target remains admissible as R increases. An assertion restricted to the single target $L = \lfloor cR^a \rfloor$ has different quantifiers and is not covered by this corollary. Only the failure at $DC - 1$ enters the threshold lower bound; the exact gap profile (I.61.10) allows other smaller interval lengths to behave differently. Replacing $L, L + 1$ by $v - 1, v$ only changes the target variable by one. At the ratio $b = 2a$ used there, (I.61.15) yields

$$a \leq \frac{1}{\nu k(k+5)/2 + 2}. \quad (\text{I.61.16})$$

For quadratics the bound is $1/(7\nu + 2)$, strengthening the earlier $1/(7\nu + 1)$ obstruction [164]. At degree one it recovers the already sharp $1/(3\nu + 2)$ threshold from [165]. For $k \geq 2$, (I.61.16) is a necessary bound; this article does not assert that it is attainable.

There is a continuous range of intermediate constraints. For $1 \leq w \leq L - 1$, the conductor is

$$C(\{L, L + 1, \dots, L + w\}) = L \left\lceil \frac{L - 1}{w} \right\rceil. \quad (\text{I.61.17})$$

To verify it, sums of exactly u permitted lengths fill the interval $[uL, u(L + w)]$. Consecutive such intervals first meet when $uw \geq L - 1$; the integer immediately preceding that first joined interval is the last gap. Substitution in (I.61.11) gives the scale

$$\mathcal{T}_{\{L, \dots, L+w\}, \mathbf{d}}(\varepsilon) \geq c'_{\mathbf{d}} \frac{L^{W+2}}{w} \prod_i \varepsilon_i^{-d_i} \quad (\text{I.61.18})$$

for a fixed positive constant $c'_{\mathbf{d}}$. When w is comparable to L , the extra tiling factor has order L ; for a fixed-width range, it has order L^2 .

Theorem I.61.4(1) also gives a lower bound of order $L^{W+1} \prod_i \varepsilon_i^{-d_i}$ for the universal problem of finding *one* good progression of length L . This has the same powers as the earlier endpoint obstruction, but is realized here by explicit rational coefficients and excludes every initial point.

I.61.5 Cyclic moduli, an exact example, and finite checks

The phases in (I.61.4) are genuine polynomial maps on $\mathbb{Z}/N\mathbb{Z}$ whenever $D \mid N$: their coefficients in that group are N/p_{ij} . For the bad interval of length $R = DC - 1$, take $N = DC = R + 1$. Thus the construction fits the interval-within-a-cyclic-group interface with an ambient modulus only one larger than the interval. This exact modulus is generally composite.

Proposition I.61.7 (Transfer of the obstruction to prime moduli). *Fix R and construct an obstruction above with tolerances $2\varepsilon_i$, where $2\varepsilon_i < 2^{-d_i}$. Choose a prime $N > 2R$ large enough that*

$$N \geq \max_i \frac{\sum_{j=1}^{d_i} (R - 1)^j}{\varepsilon_i}. \quad (\text{I.61.19})$$

For each coefficient $1/p_{ij}$ choose an integer a_{ij} nearest to N/p_{ij} . Then the polynomial phases

$$\tilde{\Phi}_i(x) = \sum_j \frac{a_{ij}}{N} x^j \pmod{1}$$

retain the obstruction at the smaller tolerances ε_i on $\{0, \dots, R - 1\}$.

Proof. At every point in this interval, the difference of the two real polynomials has absolute value at most

$$\frac{1}{2N} \sum_{j=1}^{d_i} (R-1)^j \leq \frac{\varepsilon_i}{2}.$$

If the values of $\tilde{\Phi}_i$ on a progression lie in an arc of length ε_i , expanding that arc by $\varepsilon_i/2$ at both ends contains the corresponding values of Φ_i . They therefore have diameter at most $2\varepsilon_i$. Such a progression must have step divisible by D . A prohibited partition would consequently give the same impossible residue-class tiling as before. $\square$

The extra condition $N > 2R$ also covers proper modular progressions whose carriers lie in this short integer interval. Consecutive integer representatives differ by numbers in $[-(R-1), R-1]$ that are congruent modulo N . Since this range has width less than N , those differences are equal. The progression therefore has an ordinary integer lift, with positive difference after reversal if necessary.

The prime modulus in (I.61.19) may be much larger than R . A restriction to prime ambient moduli comparable with the interval length is a different problem and is not resolved by this approximation argument.

For a small exact example, let $L = 3$, $\varepsilon = 1/16$, and

$$\Phi(x) = \frac{x}{31} + \frac{x^2}{7} \pmod{1}. \quad (\text{I.61.20})$$

The prime conditions are $31/16 < 2$ and $7/8 < 1$, so every $1/16$ -flat progression of length at least three has step divisible by $D = 217$. If the permitted lengths are three and four, their conductor is six. An interval of length $1301 = 217 \cdot 6 - 1$ has one residue class with five points; five is not a sum of threes and fours. At length 1302, every class has six points and splits into two triples of step 217. Over $\mathbb{Z}/1302\mathbb{Z}$, the same phase is represented by

$$x \mapsto 42x + 186x^2.$$

An integer arc diameter at most $\lfloor 1302/16 \rfloor = 81$ is therefore insufficient for the prohibited partition.

The accompanying standard-library checker exhaustively tests all starting residues and all step residues for three examples. It computes arc diameters with integer residues, without floating-point arithmetic.

Degree	L	Tolerance	D	Pairs tested	Good pairs
2	3	1/16	217	47,089	217
2	4	1/16	301	90,601	301
3	4	1/64	385	148,225	385

In each row the good pairs are exactly the zero step modulo D , with arbitrary starting residue. The program also checks the semigroup obstructions and a prime-modulus approximation of (I.61.20). These finite checks support the exact examples; the proofs establish the general assertions.

I.61.6 What remains to determine

The construction answers the higher-degree divisibility question in [165] and strengthens the prescribed-length obstruction from [164]. Further targets include:

1. Can general quadratic phases attain a prescribed-length upper bound at the scale $L^{3\nu+2}\varepsilon^{-2\nu}$, up to constants or subpower factors?
2. Can the leading constants be improved by replacing Lemma I.61.1 with the exact discrete minimax bound on a polynomial's leading coefficient, and by optimizing the denominator construction?
3. Does the same obstruction scale hold for prime ambient moduli restricted by $N \asymp R$?
4. Must nearly extremal obstructions exhibit a comparable hierarchy of rational coefficients, in a suitable quantified metric?
5. Can the exact gap profile (I.61.10) improve a density-increment argument that permits a controlled exceptional set?

I.62 Sources 41 and 42: Corollary 5.8 in every degree, and the quadratic variance scale [write]

Added in the eighth write, 7 October 2026.

The next twelve sections are source 41's Sections 2–7 and source 42's Sections 2–7. *Logically they close the chapter of sources 26 and 32* (Section I.40, Corollary 5.8 without a scale threshold in degrees one and two), after source 32's Section 5; they are printed here, at the end of Part I, because no number printed earlier may move. Sources 41 (Report277) and 42 (Report278) continue the series of sources 26 (Report275) and 32 (Report276). Their introductions are Sections F.46 and F.47; their source comparisons and questions are in Part V.

What source 41 proves, in one paragraph. Let $A \subseteq \mathbb{Z}/N\mathbb{Z}$ be covered by pairwise disjoint, two-comparable proper modular progressions $P_1, \dots, P_M$, with coefficient polynomials of degree at most k whose summed correlations with $\mathbf{1}_A - \delta$ are at least αN . Then some proper progression Q has $|A \cap Q| \geq (\delta + \alpha/3)|Q|$ and $|Q| \geq \max\{1, (N/M)^{1/D_k}/8\}$, $D_1 = 4$, $D_k = (k!)^2 2^{(k+1)^2}/4 + 1$, with no lower bound on α or N/M and no primality hypothesis (Theorem F.46.1). The ingredients: tent-function Fourier detection and Weyl differencing with an explicit ordered-divisor constant give $p \leq 2^{A_k} R^{B_k}$ with $\|\theta p^k\| < 1/R$ for every real θ and every integer $R \geq 2$ (Theorem I.64.3); a recursive localization keeps every point and partitions any interval of length at least $2^{c_k} L^{e_k}$ into progressions of lengths L to $2L - 1$ on which the phase varies by less than $1/(4L)$ (Theorem I.65.3), with $c_k + e_k \leq K_k/4$ (Lemma I.65.4); the subgroup-coset variance lemma (Lemma I.66.1) handles small α , and large α forces long parents (Section I.67). Section I.68 keeps the exact budgets, a compact bound $2^{-(k+8)}(N/M)^{1/(12P_k+1)}$ with limiting constants, and the sharper bounds in degrees one, two and three: the first two are source 26's Corollary I.41.3 and source 32's Theorem F.37.1, the third, $\max\{1, (N/M)^{1/4096}/8\}$, is new. The two-regime method is that of sources 26 and 32; source 41 runs it in every degree. The constants are deliberately weak; the main bound is nontrivial only for $N/M > 8^{D_k}$ (4,096 for $k = 1$, 8^{513} for $k = 2$), the variance branch may leave the cover, and no rectification is claimed.

What source 42 proves. With m the largest divisor of N below L and $q = N/m$, every $A \subseteq \mathbb{Z}/N\mathbb{Z}$ of density $0 < \delta < 1$ has a proper L -term progression of density at least $\delta + \frac{1-\delta}{L} \frac{\max\{q-L, 0\}}{q-1}$ (Theorem F.47.1), by sampling the step as a difference of independent samples from an interval of admissible steps; so the gain $(1 - \delta)/(2L)$ of the variance lemma of sources 26, 32 and 41 holds already for $N \geq 2(L - 1)^2 + 1$ in place of $N \geq L^3$, and the coefficient $1 - 1/C$ under $N \geq CL^2$ is best possible uniformly in the density, with equality on the prime-quotient sets $p\mathbb{Z}/p(L - 1)\mathbb{Z}$ (Corollary F.47.2, Theorems I.72.1, I.72.3); a sliding-window argument gives gain greater than

$\delta(1 - \delta)/(2L)$ for $N > L(L - 1)$ (Theorem I.73.1); in source 26’s affine exact-mass bound the modulus cap $(N/8)^{1/3}$ becomes $\sqrt{N/2}$ (Theorem I.74.1). Source 42 does not change source 41’s all-degree exponents (it says so), and its Proposition I.71.1 is the Kolountzakis–Révész bound.

Relation to the formal statements (checked at HEAD 6d4942222). `corollary_5_8` (Section 05.lean, line 242) asks for gain $\alpha/16$ at length $(N/M)^{1/K_k}/8$, $K_k = \text{polynomialPartitionConstant } k$ (line 45); since $D_k \leq K_k$ and $\alpha/3 \geq \alpha/16$, source 41’s theorem implies it in every degree on paper (note in Section F.46). *It is now an exact companion:* `corollary_5_8_holds` (Proofs05ThresholdFreeLocalization.lean, line 86), closed in b548d998e while this write was in progress, after the degree-one and degree-two cases `corollary_5_8_degree_one` (Proofs05AffineLocalization.lean, line 62) and `corollary_5_8_degree_two` (Proofs05QuadraticLocalization.lean, line 87). The formal proof has the architecture of this chapter: the variance branch `exists_proper_progression_density_increment` (Proofs05VarianceIncrement.lean, line 89; hypothesis $L^3 \leq N$), the high-correlation branch `density_cell_of_high_correlation` (Proofs05HighCorrelationAssembly.lean, line 53), the reduction `corollary_5_8_holds_of_threshold_free_localization` (Proofs05ThresholdFreeCorollary.lean, line 145) to the localization property `ThresholdFreePolynomialLocalization` (line 16), and that property in every degree (`thresholdFreePolynomialLocalization_all`, Proofs05ThresholdFreeLocalization.lean, line 70), proved with the repository’s own finite Weyl recurrence and budgets. Source 41’s Section I.65 proves the same property on paper with its own recurrence (note after Lemma I.65.4). Source 41’s stronger constants and source 42’s improvement of the variance scale are not formalized; the report itself gains no formal status.

I.63 An explicit quadratic recurrence estimate

Source 41, delivered Section 2.

We give the exponential-sum argument in full, including the denominator and tail bounds. All logarithms are natural.

[write] **7 October 2026; sources 40–46.** *This section is source 32’s Section 2, character for character* (Section I.45; a line diff of the two delivered sections finds no difference): Lemmas I.63.1, I.63.2, I.63.3 and Remark I.63.4 are source 32’s Lemmas I.45.1, I.45.2, I.45.3 and Remark I.45.4. It is printed again because source 41’s later sections cite these labels; the reader of source 32’s chapter may skip it.

Lemma I.63.1 (Circle approximation). *For a real β and an integer $D \geq 1$, there are integers a, q with*

$$1 \leq q \leq D, \quad \gcd(a, q) = 1, \quad \left| \beta - \frac{a}{q} \right| \leq \frac{1}{qD}.$$

There is also, without the coprimality requirement, an integer $1 \leq q \leq D$ such that $\|q\beta\| \leq 1/(D+1)$.

Proof. Consider $0, \beta, \dots, D\beta$ on the unit circle. If two coincide, their index difference gives distance zero. Otherwise some circular gap between neighboring points is at most $1/(D+1)$. The absolute difference of their indices gives $1 \leq q_0 \leq D$ and an integer a_0 with $|q_0\beta - a_0| \leq 1/(D+1)$. This also covers a gap crossing zero. Reduce a_0/q_0 to a/q . Then $q \leq q_0 \leq D$ and

$$\left| \beta - \frac{a}{q} \right| \leq \frac{1}{q_0(D+1)} \leq \frac{1}{qD}.$$

The unreduced pair supplies the second assertion. □

Lemma I.63.2 (Quadratic exponential sum). *Let $T \geq 2$ be an integer, let $\omega \in \mathbb{R}$, and put $\beta = 2\omega$. If*

$$\gcd(a, q) = 1, \quad 1 \leq q \leq T, \quad |\beta - a/q| \leq 1/(qT),$$

then

$$\left| \sum_{n=1}^T e(\omega n^2) \right|^2 \leq \frac{8T^2}{q} + 32T \log(2T). \quad (\text{I.63.1})$$

Proof. The case $q = 1$ follows from the trivial bound T^2 . Suppose $q \geq 2$ and put $b = \lfloor q/2 \rfloor \geq q/3$. Split $\{1, \dots, T-1\}$ into consecutive blocks of at most b terms. Their number is at most $T/b + 1 \leq 4T/q$.

For distinct indices u, v in the same block, $0 < |u - v| < q/2$. Coprimality and $q \leq T$ give

$$\|\beta(u - v)\| \geq \|a(u - v)/q\| - \frac{|u - v|}{qT} \geq \frac{1}{q} - \frac{1}{2T} \geq \frac{1}{2q}.$$

Thus the points βu belonging to one block have pairwise circle distance at least $1/(2q)$. Order their distances from zero as $d_1 \leq \dots \leq d_m$. An arc of length $2d$ contains at most $4qd + 1$ such points, so

$$d_j \geq \frac{j-1}{4q}.$$

Interpreting $1/0$ as infinity, it follows that

$$\begin{aligned} \sum_{u \text{ in block}} \min \left(T, \frac{1}{2\|\beta u\|} \right) &\leq T + \sum_{j=2}^m \frac{2q}{j-1} \\ &\leq T + 2q(1 + \log q). \end{aligned}$$

The last step uses $\sum_{j=1}^s 1/j \leq 1 + \log s$ for $s \geq 1$; the empty sum causes no problem. Summing over blocks gives

$$\sum_{u=1}^{T-1} \min \left(T, \frac{1}{2\|\beta u\|} \right) \leq \frac{4T^2}{q} + 8T(1 + \log q). \quad (\text{I.63.2})$$

For a geometric sum, $|1 - e(t)| = 2|\sin(\pi t)| \geq 4\|t\|$, and consequently

$$\left| \sum_{n=1}^K e(nt) \right| \leq \min \left(K, \frac{1}{2\|t\|} \right).$$

Expanding the squared quadratic sum and grouping pairs by their difference u now yields

$$\begin{aligned} \left| \sum_{n=1}^T e(\omega n^2) \right|^2 &\leq T + 2 \sum_{u=1}^{T-1} \left| \sum_{n=1}^{T-u} e(2\omega un) \right| \\ &\leq \frac{8T^2}{q} + 17T + 16T \log T \\ &\leq \frac{8T^2}{q} + 32T \log(2T). \end{aligned}$$

Here the constant phase $e(\omega u^2)$ has modulus one. The last inequality holds for $T \geq 2$, since $17 \leq 32 \log 2 + 16 \log T$. $\square$

Lemma I.63.3 (Uniform quadratic recurrence). *For every real θ and every integer $R \geq 2$, there is an integer p such that*

$$1 \leq p \leq 2^{10} R^5, \quad \|\theta p^2\| < 1/R. \quad (\text{I.63.3})$$

Proof. Set $T = 2^{10} R^5$ and $H = 4R^2$. Suppose for a contradiction that

$$\|\theta n^2\| \geq 1/R \quad (1 \leq n \leq T). \quad (\text{I.63.4})$$

On $\mathbb{R}/\mathbb{Z}$ consider the continuous triangle function

$$F(x) = \max\{1 - R\|x\|, 0\}.$$

Its Fourier coefficients are

$$c_0 = 1/R, \quad c_h = R \left(\frac{\sin(\pi h/R)}{\pi h} \right)^2 \geq 0 \quad (h \neq 0). \quad (\text{I.63.5})$$

For example, F is R times the circle convolution of the indicator of $[-1/(2R), 1/(2R)]$ with itself. Integrating that indicator against $e(-hx)$ and squaring gives (I.63.5). Also $c_h \leq R/(\pi^2 h^2)$, so the Fourier series is absolutely and uniformly convergent.

For completeness, this series indeed equals F . The Fejér kernel of order J is

$$K_J(x) = \frac{1}{J+1} \left| \sum_{v=0}^J e(vx) \right|^2.$$

It is nonnegative, its integral is one by termwise integration, and away from zero it is bounded by $1/((J+1)\sin^2(\pi x))$. Consequently convolution with K_J converges uniformly to any continuous periodic function: uniform continuity controls the near-zero integral, and the displayed bound controls the complement. Applied to F , these convolutions are also the Fejér averages of its absolutely convergent Fourier series, whose limit is the series itself. Hence $F(x) = \sum_h c_h e(hx)$, and at zero

$$\sum_{h \in \mathbb{Z}} c_h = F(0) = 1.$$

The high-frequency tail satisfies

$$\sum_{|h| > H} c_h \leq \frac{2R}{\pi^2} \sum_{h > H} h^{-2} \leq \frac{2R}{\pi^2 H} < \frac{1}{2R}. \quad (\text{I.63.6})$$

Write $S_h = \sum_{n=1}^T e(h\theta n^2)$. Assumption (I.63.4) says that $F(\theta n^2) = 0$ for all n in the sum. Interchanging a finite sum and the absolutely convergent Fourier series gives

$$0 = \frac{T}{R} + \sum_{h \neq 0} c_h S_h.$$

Using $|S_h| \leq T$ and (I.63.6), we obtain

$$\left| \sum_{0 < |h| \leq H} c_h S_h \right| > \frac{T}{2R}.$$

Since the coefficients in this sum have total at most one and $S_{-h} = \overline{S_h}$, some integer $1 \leq h \leq H$ satisfies

$$|S_h| > T/(2R). \quad (\text{I.63.7})$$

Apply Lemma I.63.1 with denominator budget T to $\beta = 2h\theta$. For the resulting relatively prime a, q , Lemma I.63.2 and (I.63.7) give

$$\frac{1}{4R^2} < \frac{8}{q} + \frac{32 \log(2T)}{T}. \quad (\text{I.63.8})$$

For $R \geq 2$, one has $\log R \leq R/2$: the difference $R/2 - \log R$ is positive at two and nondecreasing thereafter. Therefore

$$\log(2T) = 11 \log 2 + 5 \log R < 8R.$$

Indeed $11 \log 2 < 11 \leq 11R/2$. Since

$$T = 1024R^5 \geq 4096R^3 > 256R^2 \log(2T),$$

the second term on the right of (I.63.8) is less than $1/(8R^2)$. It follows that $q < 64R^2$.

Set $p = 2hq$. Then

$$1 \leq p < 512R^4 < T, \quad \|\beta q\| \leq 1/T,$$

and the exact identity $\theta p^2 = p(\beta q)$ gives

$$\|\theta p^2\| \leq p \|\beta q\| \leq \frac{p}{T} < \frac{1}{2R} < \frac{1}{R}.$$

This contradicts (I.63.4), since this integer p lies in its excluded interval. The contradiction proves (I.63.3). $\square$

Remark I.63.4. The bound $p < 512R^4$ in the last paragraph is obtained *under the contradiction hypothesis* (I.63.4). That hypothesis was used to force the large Fourier coefficient. It does not establish an unconditional R^4 recurrence theorem. The bound carried into every subsequent argument is $2^{10}R^5$.

I.64 Explicit recurrence in every degree

Source 41, delivered Section 3.

The argument is classical repeated differencing with deliberately generous constants. The point of the calculations is uniformity down to the smallest permitted scales, not a sharp recurrence exponent.

Lemma I.64.1 (An explicit ordered divisor estimate). *Let $d_s(m)$ count ordered s -tuples of positive integers with product m . For all integers $s, m \geq 1$,*

$$d_s(m) \leq 2^{3s^{2s+2}} m^{1/(2s)}.$$

Proof. For a prime power ℓ^a , unique factorization gives $d_s(\ell^a) = \binom{a+s-1}{s-1}$. This counts weak compositions of a into s parts. Each composition is realized by a word of length a over an s -letter alphabet, and words realizing different compositions are disjoint. Hence this binomial coefficient is at most s^a . If $\ell \geq s^{2s}$, it is at most $\ell^{a/(2s)}$.

For the remaining primes use $\binom{a+s-1}{s-1} \leq (a+1)^s$. Write $a = 2s^2b + t$ with $0 \leq t < 2s^2$. Since $b+1 \leq 2^b$ for every integer $b \geq 0$,

$$(a+1)^s \leq (2s^2)^s(b+1)^s \leq (2s^2)^s 2^{sb} \\ \leq (2s^2)^s 2^{a/(2s)} \leq 2^{3s^2} \ell^{a/(2s)}.$$

The final constant uses $s + 2s \log_2 s \leq 3s^2$. There are at most s^{2s} primes below s^{2s} . Multiplying the prime-power estimates therefore costs at most $2^{3s^2 s^{2s}} = 2^{3s^{2s+2}}$. For $m = 1$ the empty product is one. For $s = 1$ the small-prime range is empty, so that endpoint is included. $\square$

Lemma I.64.2 (Tracked repeated differencing). *Let $k \geq 2$, $s = k - 1$, $w = 2^s$, and $C = 2^{3s^{2s+2}}$. Let P be a real polynomial of degree at most k , with coefficient ω at x^k , and let $T \geq 2$ be an integer. Set $\beta = k!\omega$. If coprime integers a, q satisfy*

$$1 \leq q \leq T^s, \quad |\beta - a/q| \leq 1/(qT^s),$$

then

$$\left| \frac{1}{T} \sum_{n=1}^T e(P(n)) \right|^w \leq \frac{s2^{w-2}}{T} + \frac{C2^{w+1}T^{1/2}}{q} + \frac{C2^{w+2}(1 + \log q)}{T^{1/2}}. \quad (\text{I.64.1})$$

Proof. Put $S_P = \sum_{n=1}^T e(P(n))$ and $\Delta_h P(n) = P(n+h) - P(n)$. Repeated correlation expansion and Cauchy–Schwarz give, for $1 \leq j \leq s$,

$$|S_P|^{2^j} \leq (2T)^{2^j - j - 1} \sum_{|h_1|, \dots, |h_j| < T} \left| \sum_{n \in I_h} e(\Delta_{h_j} \cdots \Delta_{h_1} P(n)) \right|. \quad (\text{I.64.2})$$

Here I_h is the intersection of the original interval with the translates that keep every required shifted argument inside it; it is an integer interval, possibly empty, of length at most T . At $j = 1$ the bound is the expansion of $|S_P|^2$ followed by the triangle inequality. To pass from j to $j+1$, square the right side and use Cauchy–Schwarz over at most $(2T)^j$ shift tuples. Expand the square of each inner interval sum by correlations and take absolute values. The new prefactor exponent is

$$2(2^j - j - 1) + j = 2^{j+1} - j - 2,$$

which proves (I.64.2). Padding any shorter shift ranges to $|h| < T$ adds only nonnegative quantities.

At $j = s$, a tuple with a zero shift has identically zero phase. At most $s(2T)^{s-1}$ tuples have a zero coordinate. Bounding each inner sum by T and dividing by T^w gives their contribution $s2^{w-2}/T$.

For nonzero shifts the phase is linear, with linear coefficient $\beta h_1 \cdots h_s$. This follows by differencing the leading monomial $k-1$ times; the lower terms contribute only to the constant term. The geometric-sum bound used in Lemma I.63.2 applies to any interval of length at most T . Replacing shifts by their absolute values costs at most 2^s sign choices. For each positive product m , there are at most $d_s(m)$ positive tuples, and $m \leq T^s$. Lemma I.64.1 gives $d_s(m) \leq CT^{1/2}$. Therefore

$$|S_P/T|^w \leq \frac{s2^{w-2}}{T} + C2^{w-1}T^{-k+1/2} \sum_{m=1}^{T^s} \min \left(T, \frac{1}{2\|\beta m\|} \right). \quad (\text{I.64.3})$$

A zero distance in this minimum is interpreted as the value T . The displayed normalization is the product of $(2T)^{w-s-1}/T^w$, the 2^s sign choices, and $CT^{1/2}$; thus it contains no suppressed shift factor.

For $q \geq 2$, divide $\{1, \dots, T^s\}$ into consecutive blocks of length at most $\lfloor q/2 \rfloor$. Since $\lfloor q/2 \rfloor \geq q/3$ and $q \leq T^s$, at most $4T^s/q$ blocks suffice. In one block a nonzero index difference has $|u| < q/2$, so

$$\|\beta u\| \geq \frac{1}{q} - \frac{|u|}{qT^s} \geq \frac{1}{2q}.$$

The points in that block are $1/(2q)$ -separated on the circle. Ordering their distances from zero as $d_1 \leq \dots \leq d_b$, the interval $[-d_j, d_j]$ must accommodate j separated points, whence $d_j \geq (j-1)/(4q)$. Each block contributes at most $T + 2q(1 + \log q)$ to the sum in (I.64.3), by the harmonic-sum estimate; an empty harmonic sum is zero. Hence

$$\sum_{m=1}^{T^s} \min \left(T, \frac{1}{2\|\beta m\|} \right) \leq \frac{4T^{s+1}}{q} + 8T^s(1 + \log q).$$

For $q = 1$ this follows already from the trivial bound T^{s+1} . Substitution into (I.64.3) proves (I.64.1). $\square$

Theorem I.64.3 (Threshold free monomial recurrence). *For every integer $k \geq 2$, define*

$$\begin{aligned} s &= k-1, & w &= 2^s, & \gamma &= 3s^{2s+2}, \\ B_k &= 2w + 8 = 2^k + 8, & A_k &= 16(w + \gamma + k^2). \end{aligned} \tag{I.64.4}$$

For every real θ and every integer $R \geq 2$, there is an integer p with

$$1 \leq p \leq 2^{A_k} R^{B_k}, \quad \|\theta p^k\| < 1/R. \tag{I.64.5}$$

Proof. Set $T = 2^{A_k} R^{B_k}$, an integer, and suppose that no $1 \leq n \leq T$ has $\|\theta n^k\| < 1/R$. The tent-function proof of Lemma I.63.3 uses only the sampled real numbers, not the fact that their exponent is two. Applying exactly its Fourier coefficients (I.63.5), inversion argument, and tail estimate (I.63.6) to θn^k , some $1 \leq h \leq 4R^2$ has

$$\left| \sum_{n=1}^T e(h\theta n^k) \right| > T/(2R). \tag{I.64.6}$$

The samples are counted with multiplicity; no distinctness of their images is assumed.

Apply Lemma I.63.1 with denominator budget T^s to $\beta = k!h\theta$, and then apply Lemma I.64.2. Its left side exceeds $W = 2^{-w}R^{-w}$. We now show that the first and third right-side terms of (I.64.1) have sum less than $W/2$.

Write $A = A_k$, $B = B_k$, and $C = 2^\gamma$ in the following estimates. The first term divided by W is

$$s2^{2w-2-A}R^{-w-8} < 1/4,$$

since $s \leq 2^{k^2}$ and $A \geq k^2 + 2w$. For the third term, $q \leq T^s$, $\log R \leq R$, and $R \geq 2$ give

$$1 + \log q \leq 1 + s \log T \leq s(A + B + 1)R.$$

Its ratio to W is consequently at most

$$2^{\gamma+2w+2-A/2}s(A + B + 1)R^{-3} \leq 2^{\gamma+2w-1-A/2}s(A + B + 1).$$

Here $A \geq 64$, $B + 1 \leq A$, and $s \leq A/2$, so $s(A + B + 1) \leq A^2 \leq 2^{A/4}$. The last inequality holds at $A = 64$ and continues for larger A : the derivative of $(A/4) \log 2 - 2 \log A$ is positive there. Thus the ratio is at most

$$2^{\gamma+2w-1-A/4} = 2^{-3\gamma-2w-4k^2-1} < 1/4.$$

All these comparisons include $k = 2, R = 2$.

The middle term of (I.64.1) must therefore exceed $W/2$, yielding

$$q < C2^{2w+2}T^{1/2}R^w.$$

Set $p = k!hq$. This is a positive integer and

$$p < VT^{1/2}R^{w+2}, \quad V = k!C2^{2w+4}.$$

Since $k! \leq 2^{k^2}$,

$$\log_2 V \leq k^2 + \gamma + 2w + 4 \leq A/2;$$

the slack in the second inequality is $6w + 7\gamma + 7k^2 - 4 > 0$. Using $B/2 = w + 4$, we obtain

$$p < 2^{A/2}T^{1/2}R^{w+2} = T/R^2 < T.$$

The rational approximation gives $\|\beta q\| \leq T^{-s}$. The exact real identity $\beta q = \theta p$ and multiplication by the integer p^s imply

$$\|\theta p^k\| \leq p^s \|\beta q\| \leq (p/T)^s < R^{-2s} < 1/R.$$

This contradicts the assumed exclusion for $1 \leq p \leq T$, proving the theorem. $\square$

Remark I.64.4 (The conditional witness is not a stronger recurrence theorem). The integer with $p < T/R^2$ was manufactured after assuming that the whole interval $1 \leq n \leq T$ contains no recurrence. That assumption supplied (I.64.6). The unconditional conclusion is (I.64.5), with budget T ; the internal comparison cannot simply be promoted to a smaller global budget. Likewise, the quadratic argument carries the unconditional $2^{10}R^5$ bound, not its conditional R^4 witness.

Corollary I.64.5 (A sharper cubic constant). *For every real θ and every integer $R \geq 2$, some integer $1 \leq p \leq 2^{128}R^{16}$ satisfies $\|\theta p^3\| < 1/R$.*

Proof. For $s = 2$, $d_2(m) \leq 2^{12}m^{1/4}$. Indeed $d_2(\ell^a) = a + 1$. For primes $\ell \geq 16$, $a + 1 \leq 2^a \leq \ell^{a/4}$. For each of the six primes 2, 3, 5, 7, 11, 13, write $a = 4b + t$ with $0 \leq t < 4$; then $a + 1 \leq 4(b + 1) \leq 4 \cdot 2^b \leq 4\ell^{a/4}$. Their constants multiply to 4^6 .

Weaken this to $C = 2^{16}$ and reuse the proof of Lemma I.64.2, which at $k = 3$ gives

$$|S_P/T|^4 \leq 8/T + 32CT^{1/2}/q + 64C(1 + \log q)/T^{1/2}.$$

Take $T = 2^{128}R^{16}$ and $W = 1/(16R^4)$. Under the exclusion hypothesis, Fourier detection supplies $1 \leq h \leq 4R^2$. Apply the displayed estimate to $P(n) = h\theta n^3$, so $\beta = 6h\theta$; its left side is greater than W , and Dirichlet gives $q \leq T^2$. The first term divided by W is $2^{-121}R^{-12} < 1/4$. Moreover

$$1 + \log q \leq 1 + 2 \log T \leq 1 + 256 + 32R \leq 161R,$$

since $257 \leq 129R$ for $R \geq 2$. The third term divided by W is at most $2^{-38}161/R^3 < 1/4$. Therefore

$$q < 1024CT^{1/2}R^4, \quad p = 6hq < 24576CT^{1/2}R^6 < 2^{31}T^{1/2}R^6 < T/R^2.$$

The last comparison is $2^{95}R^{14} < 2^{128}R^{14}$. The approximation and the same integer identity give

$$\|\theta p^3\| \leq (p/T)^2 < R^{-4} < 1/R.$$

The first comparison is weak, as warranted by Dirichlet approximation; the next ones are strict. This is the desired contradiction. The general formula gives $A_3 = 3280$, $B_3 = 16$, so the sharper constant is an alternative to that general specialization. $\square$

I.65 Lossless phase partitions

Source 41, delivered Section 4.

The partitions below keep every point. A basic rounding fact is useful: a chain of length $s \geq L$ can be partitioned into chunks of lengths from L through $2L - 1$. Write $s = hL + b$ with $h \geq 1$ and $0 \leq b < L$, use $h - 1$ chunks of length L , and a last chunk of length $L + b$.

Lemma I.65.1 (Affine refinement). *Let m, L be integers with $1 \leq L \leq m$, and let $\ell, c \in \mathbb{R}$. The index interval $\{0, \dots, m - 1\}$ has a partition into arithmetic progressions C of lengths $L \leq |C| \leq 2L - 1$ and complex units z_C for which*

$$|e(-(\ell j + c)) - z_C| \leq \frac{2\pi L^2}{m} \quad (j \in C). \quad (\text{I.65.1})$$

Proof. Put $D = \lfloor m/L \rfloor \geq 1$. By the circle part of Lemma I.63.1, choose $1 \leq q \leq D$ with $\|q\ell\| \leq 1/(D + 1)$. Split the interval into its residue chains modulo q . Every chain has at least $\lfloor m/q \rfloor \geq L$ points. Apply the chunking rule above to each chain.

Write $q\ell = v + \epsilon$, where $v \in \mathbb{Z}$ and $|\epsilon| \leq 1/(D + 1)$. For a chunk $j = j_0 + tq$, $0 \leq t < s$, set

$$z_C = e(-(\ell j_0 + c))e(-\epsilon(s - 1)/2). \quad (\text{I.65.2})$$

The integer vt does not affect the exponential. The chord–arc inequality gives

$$\begin{aligned} |e(-(\ell j + c)) - z_C| &\leq 2\pi|\epsilon| \left| t - \frac{s - 1}{2} \right| \\ &\leq \pi|\epsilon|(s - 1) \leq \frac{2\pi(L - 1)}{D + 1} < \frac{2\pi L^2}{m} \end{aligned}$$

when $L > 1$, using $D + 1 > m/L$. When $L = 1$, every chunk is a singleton and the error is zero. This proves the stated weak bound in all cases. $\square$

The reduction before taking the midpoint in (I.65.2) matters. If $q\ell = 1$ and $s = 2$, both sampled phase values agree, while the midpoint obtained from the unreduced increment contributes $e(-1/2) = -1$. The reduced increment is zero and its midpoint has the correct phase. Integer shifts are harmless at integer parameters but need not be harmless at half-integer midpoints.

Lemma I.65.2 (Quadratic refinement for real coefficients). *Let n, L be integers with $L \geq 2$ and $n \geq 2^{116}L^{38}$. For every real polynomial p_0 of degree at most two on $\{0, \dots, n - 1\}$ there is a partition into integer arithmetic progressions C , with units z_C , such that*

$$L \leq |C| \leq 2L - 1, \quad |e(-p_0(j)) - z_C| < 1/(4L) \quad (j \in C). \quad (\text{I.65.3})$$

Proof. Write $p_0(j) = \theta j^2 + \beta j + \gamma_0$ with real coefficients, and set

$$H = 64L^3 = 2^6L^3, \quad R = 256LH^2 = 2^{20}L^7, \quad T = 2^{10}R^5 = 2^{110}L^{35}. \quad (\text{I.65.4})$$

All are integers, and $R \geq 2$. Lemma I.63.3 supplies $1 \leq p \leq T$ with $\|\theta p^2\| < 1/R$. Since $n \geq TH \geq pH$, each residue chain modulo p in the index interval has at least H terms. Chunk each into lengths H through $2H - 1$ by the rule above.

On a piece $j = j_0 + pt$, $0 \leq t < m$, where $H \leq m \leq 2H - 1$, expansion gives

$$p_0(j_0 + pt) = \theta p^2 t^2 + \ell t + c_0, \quad \ell = 2\theta j_0 p + \beta p, \quad c_0 = p_0(j_0).$$

Choose $v \in \mathbb{Z}$ with $\epsilon = \theta p^2 - v$ and $|\epsilon| < 1/R$. Subtracting vt^2 leaves the exponential unchanged at integer t . The chord-arc inequality gives

$$|e(-p_0(j_0 + pt)) - e(-(\ell t + c_0))| \leq 2\pi|\epsilon|t^2 < 8\pi H^2/R = \pi/(32L).$$

Apply Lemma I.65.1 to the remaining affine phase at target L . Its error is at most $2\pi L^2/m \leq 2\pi L^2/H = \pi/(32L)$. Their sum is less than $\pi/(16L) < 1/(4L)$. Each final cell $t = t_0 + qv$ gives an arithmetic progression $j = j_0 + pt_0 + pqv$ in the original interval. These cells remain disjoint and retain every index. Finally $TH = 2^{116}L^{38}$, as required. $\square$

Theorem I.65.3 (Recursive polynomial localization). *Set $c_2 = 116$, $e_2 = 38$. For each integer $k \geq 3$, use A_k, B_k from (I.64.4) and define*

$$\begin{aligned} m_k &= kB_k + 1, \\ e_k &= B_k + m_k e_{k-1}, \\ c_k &= A_k + (k+7)B_k + m_k(c_{k-1} + e_{k-1}), \\ F_k(L) &= 2^{c_k} L^{e_k}. \end{aligned} \quad (\text{I.65.5})$$

Also set $F_2(L) = 2^{116}L^{38}$. For integers n, L with $L \geq 2$ and $n \geq F_k(L)$, every real polynomial p_0 of degree at most k on $\{0, \dots, n-1\}$ has a partition into arithmetic progressions C , with units z_C , such that

$$L \leq |C| \leq 2L - 1, \quad |e(-p_0(j)) - z_C| < 1/(4L) \quad (j \in C). \quad (\text{I.65.6})$$

For degree at most one, $n \geq 64L^3$ is sufficient.

Proof. The degree-two assertion is Lemma I.65.2. For degree one, Lemma I.65.1 gives error at most $2\pi L^2/n \leq \pi/(32L) < 1/(4L)$.

Suppose $k \geq 3$ and the assertion holds at degree $k-1$. Set

$$H = F_{k-1}(2L) = 2^{c_{k-1}+e_{k-1}} L^{e_{k-1}}, \quad R = 2^{k+7} L H^k, \quad T = 2^{A_k} R^{B_k}. \quad (\text{I.65.7})$$

Every scale is an integer and $R \geq 2$. The definitions give the exact identity

$$TH = 2^{A_k+(k+7)B_k} L^{B_k} H^{kB_k+1} = 2^{c_k} L^{e_k} = F_k(L). \quad (\text{I.65.8})$$

Let θ be the coefficient of x^k in p_0 . Then Theorem I.64.3 supplies $1 \leq p \leq T$ with $\|\theta p^k\| < 1/R$. This coefficient is zero when the degree is below k . Split the index interval into residue chains modulo p . Each has at least H terms since $n \geq TH \geq pH$. The chunking rule splits each chain into pieces of lengths H through $2H - 1$.

On one such piece write $j = j_0 + pt$, $0 \leq t < m$, with $H \leq m \leq 2H - 1$. For some real polynomial ψ of degree at most $k - 1$ and integer v ,

$$p_0(j_0 + pt) = vt^k + \epsilon t^k + \psi(t), \quad |\epsilon| < 1/R.$$

The coefficient at t^k before reduction is exactly θp^k . At integer t the term vt^k disappears from the exponential, so

$$|e(-p_0(j_0 + pt)) - e(-\psi(t))| \leq 2\pi|\epsilon|t^k < \frac{2\pi(2H)^k}{R} = \frac{\pi}{64L}.$$

Apply the induction hypothesis to ψ at target length $2L$. This is legal because $m \geq H = F_{k-1}(2L)$. Its cells have lengths $2L$ through $4L - 1$, with constant-unit error less than $1/(8L)$. Split each such cell again by the chunking rule into lengths L through $2L - 1$. Restrict its same unit to the smaller cells; this does not increase the error. The combined error is

$$< \frac{\pi}{64L} + \frac{1}{8L} < \frac{3}{16L} < \frac{1}{4L}.$$

Every stage partitions its input exactly and composes arithmetic parameterizations, so all final cells are arithmetic progressions and every original index occurs once. $\square$

Lemma I.65.4 (The uniform source budget). *For all integers $k \geq 2$,*

$$c_k + e_k \leq K_k/4. \tag{I.65.9}$$

Consequently the sufficient condition $n \geq (2L)^{K_k/4}$ implies (I.65.6) for integer $L \geq 2$.

Proof. Let $U_k = c_k + e_k$. The base is $U_2 = 154 \leq 2048/4$. For $k \geq 3$, (I.65.5) gives

$$U_k = A_k + (k + 8)B_k + m_k(c_{k-1} + 2e_{k-1}) \leq A_k + (k + 8)B_k + 2m_kU_{k-1}. \tag{I.65.10}$$

We bound the three terms separately. First $k - 1 \leq 2^{k/2}$ for $k \geq 2$. Check $k = 2, 3, 4$ directly; thereafter the induction follows from $k/(k-1) \leq 4/3 < \sqrt{2}$. Thus $(k-1)^{2k} \leq 2^{k^2}$; also $2^{k-1} + k^2 \leq 2^{k^2}$ for $k \geq 3$, since each summand is at most 2^{k^2-1} . Therefore

$$A_k \leq 64 \cdot 2^{k^2}, \quad \frac{A_k}{K_k} \leq \frac{64}{(k!)2^{2k+1}} \leq \frac{64}{36 \cdot 128} = \frac{1}{72} < \frac{1}{64}.$$

Second, $B_k \leq 2^{k+1}$ and $k + 8 \leq 4k$ for $k \geq 3$. Since $k! \geq k$ and $(k+1)^2 \geq k+9$,

$$(k + 8)B_k \leq k2^{k+3} \leq K_k/64.$$

Third, $m_k \leq k2^{k+2}$ and

$$K_k/K_{k-1} = k^2 2^{2k+1}, \quad \frac{m_k}{K_k/K_{k-1}} \leq \frac{1}{k2^{k-1}} \leq \frac{1}{12}.$$

By the induction hypothesis, $2m_kU_{k-1} \leq m_kK_{k-1}/2 \leq K_k/24$. Substituting into (I.65.10) yields

$$U_k \leq \left(\frac{1}{64} + \frac{1}{64} + \frac{1}{24} \right) K_k = \frac{7}{96} K_k < K_k/4.$$

This is an induction for every integer k , with no finite range or asymptotic cutoff left unchecked. Finally

$$F_k(L) = 2^{c_k} L^{e_k} \leq (2L)^{c_k + e_k} \leq (2L)^{K_k/4},$$

which proves the sufficient condition. $\square$

[write] 7 October 2026; sources 40–46. *It supplies the formal localization hypothesis, on paper* (checked at the write). The repository’s `ThresholdFreePolynomialLocalization k` (`Proofs05ThresholdFreeCorollary.lean`, line 16, at HEAD 6d4942222) asks: for every proper modular progression P with $(2L)^{\lfloor K_k/2 \rfloor} \leq |P|$, $L \geq 2$, and every ϕ with `PolynomialOn k`, a partition of P into proper progressions of length at least L with units z_j and $\|e(-\phi(x)) - z_j\| \leq 1/(4L)$ on each. For $k \geq 2$, Theorem I.65.3 with this lemma gives cells of lengths L to $2L - 1$ and error $< 1/(4L)$ once $|P| \geq (2L)^{K_k/4}$, which is at most $(2L)^{K_k/2}$ (K_k is even; $c_k + e_k \leq K_k/4$ was confirmed in exact integers for $k \leq 60$ at the write); the passage to modular parents is the paragraph after the lemma. For $k = 1$, Lemma I.65.1 gives error at most $2\pi L^2/n \leq \pi/(32L)$ once $n \geq 64L^3$, and $64L^3 \leq (2L)^8$ for $L \geq 2$. The repository’s `corollary_5_8_holds_of_threshold_free_localization` (line 145) derives `corollary_5_8` from this hypothesis for every $k \geq 1$, and the repository has since proved the hypothesis itself (`thresholdFreePolynomialLocalization_all`, `Proofs05ThresholdFreeLocalization.lean`, line 70) with its own finite Weyl recurrence, radius exponent $2((k-1)^2 + 1)(2^{k-1} + 3)$ in place of this section’s $B_k = 2^k + 8$; so this section is a second, written route to the formal hypothesis, not the formalized one.

Passage to modular parents. If $P = \{u + jd : 0 \leq j < n\}$ is proper and $\phi(x) = \sum_{v=0}^k a_v x^v$ over G , choose integer representatives $\tilde{a}_v, \tilde{u}, \tilde{d}$ and set

$$p_0(j) = \frac{1}{N} \sum_{v=0}^k \tilde{a}_v (\tilde{u} + j\tilde{d})^v.$$

This is a real coefficient polynomial of degree at most k and $e_N(-\phi(u + jd)) = e(-p_0(j))$ for every integer j . Integer representatives may change the normalized value by an integer, leaving the phase unchanged even through modular wrapping. The real-variable localization therefore applies. Each index cell maps to a modular progression, and injectivity of the parent map restricts to injectivity on the cell. Cardinalities and disjointness are preserved. No division in G and no invertibility of d or $k!$ is used.

Two cubic budgets. The general recursion has $A_3 = 3280$, $B_3 = 16$, $m_3 = 49$, hence $(c_3, e_3) = (10986, 1878)$. Replacing only its cubic recurrence input by Corollary I.64.5 gives instead

$$c'_3 = 128 + 10 \cdot 16 + 49(116 + 38) = 7834, \quad e'_3 = 16 + 49 \cdot 38 = 1878. \quad (\text{I.65.11})$$

The same localization proof gives $n \geq 2^{7834} L^{1878}$. The all-degree induction above uses the unmodified constants; the sharper cubic option does not alter that proof.

I.66 A second moment over every cyclic modulus

Source 41, delivered Section 5.

Over a prime modulus, pairs of distinct progression positions often have a particularly simple distribution. Over an arbitrary cyclic modulus their distribution is instead controlled by additive subgroup cosets. The resulting inequality is enough here.

[write] 7 October 2026; sources 40–46. *This section is source 32’s Section 4, character for character* (Section I.47); Lemma I.66.1 is source 32’s Lemma I.47.1 and source 26’s Lemma I.42.1. *Formalized since the seventh write* (checked at HEAD 6d4942222): `exists_proper_progression_density_increment` (`Proofs05VarianceIncrement.lean`, line 89) proves exactly this statement for $A \subseteq \mathbb{Z}/N\mathbb{Z}$, $2 \leq L$, $L^3 \leq N$, $0 < \delta < 1$, a proper progression of length L with density at least $\delta + (1 - \delta)/(2L)$; its docstring credits Report277. The ingredients are `cyclic_progression_second_moment_lower` (`Proofs05ProgressionMoments.lean`,

line 74) and `badProgressionDirections_card_le` (`Proofs05ProperDirections.lean`, line 14). Source 42 proves the same gain under $N \geq 2(L-1)^2 + 1$ (Corollary F.47.2); that version is not formalized.

Lemma I.66.1 (Proper progression from variance). *Let $0 < \delta < 1$, let $L \geq 2$ be an integer, and suppose $N \geq L^3$. There is a proper L -term modular progression Q with*

$$\text{dens}_Q(A) := \frac{|A \cap Q|}{|Q|} \geq \delta + \frac{1 - \delta}{2L}. \quad (\text{I.66.1})$$

Proof. Choose a, d independently and uniformly in G , including $d = 0$, and put

$$X(a, d) = \frac{1}{L} \sum_{j=0}^{L-1} \mathbf{1}_A(a + jd).$$

Repeated positions are counted with their index multiplicity. For every fixed d , averaging over a gives $\mathbb{E}_a X = \delta$.

For distinct j, k , let $H = (k - j)G$. The change of variables $(a, d) \mapsto (x, d) = (a + jd, d)$ is a bijection of G^2 . The element $(k - j)d$ is uniform in H , independently of x . Therefore

$$\mathbb{E}[\mathbf{1}_A(a + jd)\mathbf{1}_A(a + kd)] = \sum_{C \in G/H} \frac{|H|}{N} \left(\frac{|A \cap C|}{|H|} \right)^2 \geq \delta^2. \quad (\text{I.66.2})$$

Indeed the coset densities have weighted mean δ , and their weighted squared deviations from that mean sum to a nonnegative number. Thus all off-diagonal covariances are nonnegative, and the L diagonal terms yield

$$\text{Var}_{a,d} X \geq \frac{\delta(1 - \delta)}{L}. \quad (\text{I.66.3})$$

A direction d is bad precisely when the L -term parameterization is not injective. This means $hd = 0$ for some $1 \leq h < L$. The kernel of multiplication by h on G has $\gcd(h, N) \leq h$ elements: if $g = \gcd(h, N)$, then $N \mid hd$ is equivalent to $N/g \mid d$. Hence the proportion η of bad directions obeys

$$\eta \leq \frac{1}{N} \sum_{h=1}^{L-1} h = \frac{L(L-1)}{2N} < \frac{1}{2L}. \quad (\text{I.66.4})$$

At least one proper direction exists, since $d = 1$ works for $L \leq N$.

Let h_* be the largest value of $X(a, d) - \delta$ over proper directions and all a . For a proper direction, $0 \leq X \leq \delta + h_*$ implies $X^2 \leq (\delta + h_*)X$, and so $\text{Var}_a X \leq \delta h_*$. For any bad direction, $0 \leq X \leq 1$ similarly gives $\text{Var}_a X \leq \delta(1 - \delta)$. Since every conditional mean is δ , averaging these variances and using (I.66.3) gives

$$\frac{\delta(1 - \delta)}{L} \leq (1 - \eta)\delta h_* + \eta\delta(1 - \delta).$$

Dividing by $\delta > 0$ and rearranging,

$$h_* \geq (1 - \delta) \frac{1/L - \eta}{1 - \eta} \geq \frac{1 - \delta}{2L},$$

where (I.66.4) and $1 - \eta \leq 1$ justify the last step. A maximizing proper progression proves the claim. $\square$

I.67 Assembly at every scale

Source 41, delivered Section 6.

We first collect the cover algebra, since the same quantities will also govern the sharper alternatives. Put

$$b = 1 - \delta, \quad \tau = |U|/N, \quad r = \min_i |P_i|, \quad S = N/M, \quad \mu = \frac{1}{N} \sum_{x \in U} |f(x)|.$$

Nonempty disjointness gives $M \leq N$ and $S \geq 1$; coverage gives $\delta \leq \tau \leq 1$. Counting points of A and $U \setminus A$ separately gives

$$\mu = \delta(1 + \tau - 2\delta), \quad 0 < \alpha \leq \mu \leq 2\delta b \leq 2\tau b \leq 2b. \quad (\text{I.67.1})$$

Here the correlation and triangle inequality imply $\alpha \leq \mu$, which rules out $\delta = 0, 1$. Thus $0 < \delta < 1$, $b > 0$, and A is nonempty. Comparability yields

$$\tau N = \sum_i |P_i| \leq 2Mr, \quad r \geq \tau S/2. \quad (\text{I.67.2})$$

Lemma I.67.1 (From a full phase partition to a positive cell). *Suppose the parents are partitioned into proper cells C , with units z_C . Let $i(C)$ be the parent of C , put $F_C = \sum_{x \in C} f(x) \in \mathbb{R}$, and define*

$$E = \sum_C \sum_{x \in C} |f(x)| |e_N(-\phi_{i(C)}(x)) - z_C|.$$

Then

$$\sum_C |F_C| \geq \alpha N - E, \quad \sum_C F_C = \delta(N - |U|) \geq 0. \quad (\text{I.67.3})$$

If $E \leq \alpha N/3$, some cell has density at least $\delta + \alpha/3$. If the error inequality is strict, so is the resulting density inequality.

Proof. For each parent, replace the phase by z_C on each cell, bound the replacement error, and apply the triangle inequality. Since $|z_C| = 1$, summing gives the first inequality in (I.67.3). Every point of U occurs exactly once, and $A \subseteq U$, so the signed sum is $|A| - \delta|U| = \delta(N - |U|)$. The positive F_C sum to

$$\frac{1}{2} \left(\sum_C |F_C| + \sum_C F_C \right) \geq \frac{\alpha N - E}{2} \geq \frac{\alpha N}{3}.$$

These cells have total cardinality at most N . Their weighted average discrepancy is therefore at least $\alpha/3$, and some cell has at least that discrepancy. The same calculation preserves strictness. Coverage gives a nonnegative signed total, not necessarily zero. $\square$

Proof of Theorem F.46.1. Set $x = S^{1/D_k}/8$ and $L = \lceil x \rceil$. If $L = 1$, take a singleton in A . Its length is at least x , and its density is one, strictly larger than $\delta + \alpha/3$, since (I.67.1) gives $\alpha/3 \leq 2\delta b/3 < b$.

Suppose $L \geq 2$. Then $x > 1$ and $L < x + 1 < 2x$, whence

$$S > (4L)^{D_k} > L^3, \quad N \geq S > L^3. \quad (\text{I.67.4})$$

If $\alpha \leq 3b/(2L)$, Lemma I.66.1 supplies a proper L -term progression of density at least $\delta + b/(2L) \geq \delta + \alpha/3$. Its length is at least x .

Otherwise $\alpha > 3b/(2L)$. Together with (I.67.1) and (I.67.2), this implies

$$\tau > 3/(4L), \quad r > 3S/(8L). \quad (\text{I.67.5})$$

For $k = 1$, $D_1 = 4$, so (I.67.4) gives $r > 96L^3 > 64L^3$. For $k \geq 2$, put $t = K_k/4$, so $D_k = t + 1$; then

$$r > \frac{3}{8} \frac{(4L)^{t+1}}{L} = \frac{3}{2} 4^t L^t > (2L)^t.$$

Thus every parent meets the affine or all-degree localization threshold. Apply Theorem I.65.3 and Lemma I.65.4, with the modular passage established there, at the same target length L . The total phase error satisfies

$$E < \frac{\mu N}{4L} \leq \frac{bN}{2L} < \frac{\alpha N}{3}.$$

Strictness in the first comparison follows from the strict pointwise error and $\mu > 0$. Lemma I.67.1 supplies a proper cell of density greater than $\delta + \alpha/3$ and length at least $L \geq x$. Every branch gives a nonempty progression, which proves the maximum with one in (F.46.4). $\square$

The high-correlation branch finds Q inside a parent. The variance branch searches all directions in the whole group. Requiring parent containment in the final theorem would add a conclusion not established by that branch.

I.68 Exact budgets and sharper alternatives

Source 41, delivered Section 7.

The uniform K_k comparison is useful for matching the source interface. It intentionally loses information. Neither its degree-one nor its degree-two specialization replaces the stronger small-degree bounds below.

Corollary I.68.1 (The actual recursive budget). *Under the cover and correlation hypotheses of Theorem F.46.1, suppose that degree at most k phases have a valid partition threshold $n \geq 2^c L^e$ for integer $L \geq 2$, giving cells of lengths L through $2L - 1$ and phase error less than $1/(4L)$. Here $c \geq 0$, $e \geq 2$ are fixed real numbers. Put*

$$R_0 = \left(\frac{3S}{8 \cdot 2^c} \right)^{1/(e+1)}, \quad L_0 = \max\{1, \lfloor R_0 \rfloor\}.$$

There is a proper progression Q with

$$|Q| \geq L_0 \geq \max\{1, R_0/2\}, \quad |A \cap Q| \geq (\delta + \alpha/3)|Q|. \quad (\text{I.68.1})$$

Available choices are $(c, e) = (6, 3)$ for degree one, $(116, 38)$ for degree two, (c_k, e_k) from (I.65.5) for every $k \geq 3$, and $(7834, 1878)$ for the sharper cubic alternative.

Proof. If $L_0 = 1$, a singleton in A works as before. Otherwise $L_0 \leq R_0$, and

$$S \geq \frac{8 \cdot 2^c}{3} L_0^{e+1} \geq L_0^3, \quad N \geq S.$$

If $\alpha \leq 3b/(2L_0)$, apply Lemma I.66.1. Otherwise (I.67.5), with $L = L_0$, gives

$$r > 3S/(8L_0) \geq 2^c L_0^e.$$

Every parent therefore meets the stipulated partition threshold, and its total phase error is less than $\mu N/(4L_0) \leq bN/(2L_0) < \alpha N/3$. Lemma I.67.1 finishes this branch. If $R_0 < 2$, the last lower bound in (I.68.1) is one; if $R_0 \geq 2$, then $\lfloor R_0 \rfloor \geq R_0/2$. This includes all floor boundaries without an unproved real-root rounding step. $\square$

For fixed k , the actual power of S in this corollary is $1/(e_k + 1)$, accompanied by its explicit constant. In particular the sharper cubic budget gives a power $1/1879$ with a substantial fixed prefactor. One should compare the complete guarantees, including the singleton floor and constants, rather than compare exponents alone.

Corollary I.68.2 (A compact bound for the degree dependence). *For $k \geq 2$, put $\mathcal{P}_k = k! 2^{k(k+1)/2}$. The unmodified recursive budgets in (I.65.5) satisfy*

$$e_k < 12\mathcal{P}_k, \quad c_k < (k+7)e_k, \quad c_k + e_k < 12(k+8)\mathcal{P}_k. \quad (\text{I.68.2})$$

Consequently, for integers $n, L \geq 2$, the condition $n \geq (2^{k+7}L)^{12\mathcal{P}_k}$ suffices for the degree- k localization at target L . Under Theorem F.46.1's hypotheses, a proper progression has gain $\alpha/3$ and length at least

$$\max \left\{ 1, 2^{-(k+8)} S^{1/(12\mathcal{P}_k+1)} \right\}. \quad (\text{I.68.3})$$

Here $\mathcal{P}_k = 2^{-(k+1)/2} \sqrt{K_k}$. This is a bound on the growth of the proved sufficient budget, not an optimality statement. The calligraphic symbol distinguishes this numerical quantity from a parent progression P_i .

Proof. Let $\Pi_2 = 1$ and $\Pi_k = \prod_{j=3}^k m_j$. Dividing the e_k recurrence by Π_k gives

$$E_k := \frac{e_k}{\Pi_k} = 38 + \sum_{j=3}^k \frac{B_j}{\Pi_j}.$$

The m_j increase, $m_3 = 49$, $m_4 = 97$, and $B_j/m_j < 1/j$. Let $E_\infty = 38 + \sum_{j=3}^\infty B_j/\Pi_j$; the following comparison proves convergence. Since $\Pi_{j-1} \geq 49 \cdot 97^{j-4}$ for $j \geq 4$,

$$E_k \leq E_\infty < 38 + \frac{16}{49} + \frac{1}{4 \cdot 49} \sum_{t=0}^\infty 97^{-t} = 38 + \frac{16}{49} + \frac{97}{18816} < \frac{115}{3}. \quad (\text{I.68.4})$$

The final comparison uses $97/18816 < 1/147$.

Set $d_j = (8 + 1/j)/2^j$ and $q_k = \Pi_k/\mathcal{P}_k$. Then

$$q_k = \frac{1}{16} \prod_{j=3}^k (1 + d_j), \quad q_7 = \frac{18772640957}{64424509440} < \frac{7}{24}.$$

The last strict comparison is $18772640957 < 18790481920$. Also $\sum_{j=8}^\infty d_j \leq (8 + 1/8) \sum_{j=8}^\infty 2^{-j} = 65/1024$. For nonnegative x_i of sum at most $t < 1$, expanding a finite product and bounding each elementary symmetric sum by t^r gives $\prod_i (1 + x_i) \leq \sum_{r \geq 0} t^r = 1/(1 - t)$. This bounds the increasing finite products and proves convergence here. Write $q_\infty = \lim_{k \rightarrow \infty} q_k$. Consequently, including $k < 7$ by monotonicity,

$$q_k \leq q_\infty < \frac{7/24}{1 - 65/1024} = \frac{896}{2877}, \quad \frac{e_k}{\mathcal{P}_k} < \frac{896}{2877} \frac{115}{3} = \frac{103040}{8631} < 12.$$

For the coefficient bound let $J_k = (k + 7)e_k - c_k$. Direct substitution into (I.65.5) yields

$$J_2 = 226, \quad J_k = m_k J_{k-1} - A_k, \quad \frac{J_k}{\Pi_k} = 226 - \sum_{j=3}^k a_j, \quad a_j = \frac{A_j}{\Pi_j}. \quad (\text{I.68.5})$$

We show that the infinite sum is less than 226. For $j \geq 9$,

$$a_{j+1}/a_j < 1/8. \quad (\text{I.68.6})$$

Indeed, compare the three positive summands 2^{j-1} , $3(j-1)^{2j}$, and j^2 defining $A_j/16$, and divide their respective next-term ratios by m_{j+1} . The first ratio is $2/m_{j+1} < 1/8$. The third is $((j+1)/j)^2/m_{j+1} < 2/m_{j+1} < 1/8$. The middle is

$$\frac{j^2(1 + 1/(j-1))^{2j}}{m_{j+1}}.$$

For $n \geq 1$, the binomial expansion gives $(1 + 1/n)^n \leq \sum_{r=0}^n 1/r! < 3$; here $r! \geq 2^{r-1}$ for $r \geq 2$, with strictness in the finite sum. Therefore, for $j \geq 9$,

$$(1 + 1/(j-1))^{2j} < 9(j/(j-1))^2 \leq 9(9/8)^2 < 12.$$

Since $m_{j+1} \geq (j+1)2^{j+1}$, the middle ratio is less than $12j/2^{j+1} \leq 108/1024 = 27/256 < 1/8$. The sequence $j/2^{j+1}$ is nonincreasing for $j \geq 1$. Each positive summand has the required ratio, proving (I.68.6) for their sum.

For the finite calculation, the following integers z_j satisfy the exact certificates

$$z_j \Pi_j \leq 100000 A_j < (z_j + 1) \Pi_j. \quad (\text{I.68.7})$$

j	3	4	5	6	7	8	9	10
z_j	6693877	6633957	5268450	2832889	954144	191498	22175	1450

Each certificate is integer substitution into the displayed formulas for A_j and Π_j . They are also reproduced by the exact companion. Summing the upper bounds through $j = 9$, and retaining the $j = 10$ bound for the tail, gives

$$\sum_{j=3}^9 a_j < \frac{22596997}{100000} < \frac{22597}{100}, \quad a_{10} < \frac{1451}{100000} < \frac{3}{200}.$$

Thus (I.68.6) implies

$$\sum_{j=3}^{\infty} a_j < \frac{22597}{100} + \frac{8}{7} \frac{3}{200} = 226 - \frac{9}{700} < 226.$$

Equation (I.68.5) now gives $J_k > 0$ for every $k \geq 2$. This proves (I.68.2) with no unverified degree range.

The partition conclusion follows from $2^{c_k} L^{e_k} < (2^{k+7} L)^{12\mathcal{P}_k}$. Apply Corollary I.68.1 with integer exponents $C = 12(k+7)\mathcal{P}_k$ and $E = 12\mathcal{P}_k$. Its real-valued length term is $R_0/2$, and the exact identity

$$\frac{R_0}{2} = 2^{-(k+8)} (3 \cdot 2^{k+4} S)^{1/(12\mathcal{P}_k+1)} \geq 2^{-(k+8)} S^{1/(12\mathcal{P}_k+1)}$$

proves (I.68.3). Finally, the comparison with K_k is direct algebra from their definitions. $\square$

Corollary I.68.3 (Exact limiting constants). *Retain Π_k, q_k, E_k, a_j from the proof of Corollary I.68.2, and define*

$$\Theta_\infty = 226 - \sum_{j=3}^{\infty} a_j, \quad \lambda = q_\infty E_\infty, \quad \gamma_* = 7 - \Theta_\infty / E_\infty.$$

These are constants defined by convergent series and a convergent product, and

$$11 < \lambda < 12, \quad 7 - \frac{1}{1900} < \gamma_* < 7 - \frac{27}{80500} < 7. \quad (\text{I.68.8})$$

As $k \rightarrow \infty$ through integers,

$$\frac{e_k}{\mathcal{P}_k} = \lambda + O(2^{-k}), \quad \frac{c_k}{e_k} - k \rightarrow \gamma_*, \quad \frac{c_k}{\mathcal{P}_k} = \lambda(k + \gamma_*) + O(k2^{-k}). \quad (\text{I.68.9})$$

In particular $c_k/\mathcal{P}_k - \lambda k \rightarrow \lambda\gamma_$. The constants and error terms describe this proved construction, not the optimal density-increment problem.*

Proof. The previous proof gives $\Theta_\infty > 9/700$. The lower integer certificates (I.68.7), summed through $j = 10$, give

$$\sum_{j=3}^{10} a_j \geq \frac{22598440}{100000} > \frac{11299}{50} = 226 - \frac{1}{50},$$

so $\Theta_\infty < 1/50$. Since $38 \leq E_\infty < 115/3$, the bounds on γ_* follow. Also $\lambda < 12$, while $\lambda > 38q_7 > 11$, the latter a direct integer comparison using the displayed value of q_7 .

Here are explicit rational errors, which also justify the rates in (I.68.9). For $k \geq 4$ put

$$s_k = \frac{8 + 1/(k+1)}{2^k}, \quad t_k = \frac{m_{k+1}}{(k+1)(m_{k+1} - 1)\Pi_k}.$$

Then $s_k < 1$. The same product estimate as before and the tail identity for E_k give

$$0 < E_\infty - E_k < t_k, \quad q_k < q_\infty \leq \frac{q_k}{1 - s_k}. \quad (\text{I.68.10})$$

Indeed $\sum_{j>k} d_j \leq s_k$, and for $j \geq k+1$ one has $B_j/m_j < 1/j \leq 1/(k+1)$ and $\Pi_{j-1} \geq \Pi_k m_{k+1}^{j-k-1}$; summing the resulting geometric series gives t_k . In particular, splitting the difference $q_\infty E_\infty - q_k E_k$ yields

$$0 < \lambda - \frac{e_k}{\mathcal{P}_k} < 12s_k + q_k t_k = 12s_k + \frac{m_{k+1}}{(k+1)(m_{k+1} - 1)\mathcal{P}_k}. \quad (\text{I.68.11})$$

We used $q_\infty - q_k \leq s_k q_\infty$ and $q_\infty E_\infty < 12$.

For $k \geq 9$, put $h_k = (8/7)a_{k+1}$, $\Theta_k = 226 - \sum_{j=3}^k a_j$, and $g_k = c_k/e_k - k$. Equation (I.68.5) gives $g_k = 7 - \Theta_k/E_k$. The geometric tail (I.68.6) implies $0 < \Theta_k - \Theta_\infty < h_k$, so

$$\begin{aligned} 0 < \gamma_* - g_k &= \frac{\Theta_k - \Theta_\infty}{E_k} + \frac{\Theta_\infty(E_\infty - E_k)}{E_k E_\infty} \\ &< \frac{h_k}{38} + \frac{t_k}{50 \cdot 38^2}. \end{aligned} \quad (\text{I.68.12})$$

Finally $c_k/\mathcal{P}_k = (e_k/\mathcal{P}_k)(k + g_k)$, and hence

$$0 < \lambda(k + \gamma_*) - \frac{c_k}{\mathcal{P}_k} < (k+7)(12s_k + q_k t_k) + 12 \left(\frac{h_k}{38} + \frac{t_k}{50 \cdot 38^2} \right). \quad (\text{I.68.13})$$

Now $s_k = O(2^{-k})$, $q_k t_k = O(1/(k\mathcal{P}_k))$, $\Pi_k \geq 49^{k-2}$, and (I.68.6) gives $h_k = O(8^{-k})$. These explicit bounds prove all limits and rates claimed; no decimal approximation of an infinite sum is needed. $\square$

Corollary I.68.4 (Retained affine bounds). *When the phases have degree at most one, the same gain $\alpha/3$ is available with each of the length guarantees*

$$\min \left\{ \left(\frac{rb}{128\mu} \right)^{1/3}, (N/8)^{1/3} \right\}, \quad \sqrt{\frac{\alpha r}{24\pi\mu}}. \quad (\text{I.68.14})$$

The second progression can be chosen inside one parent. For the two-comparable cover in Theorem F.46.1, one may consequently choose a progression of length at least

$$\max \left\{ 1, \frac{S^{1/3}}{8}, \sqrt{\frac{\alpha S}{96\pi b}} \right\}. \quad (\text{I.68.15})$$

Parent containment is assured for the square-root alternative, not for an arbitrary selection attaining the maximum.

Proof. For any integer $1 \leq L \leq r$, apply Lemma I.65.1 to each real pullback phase. The resulting error in Lemma I.67.1 is at most

$$E \leq 2\pi L^2 \mu N / r. \quad (\text{I.68.16})$$

For the cube-root guarantee, put $H_* = rb/\mu$, choose $x = \min\{(H_*/128)^{1/3}, (N/8)^{1/3}\}$ and $L = \lceil x \rceil$. A singleton handles $L = 1$. Otherwise $L < 2x$, hence $H_* > 16L^3$ and $N > L^3$. If $\alpha \leq 3b/(2L)$, use the variance lemma. In the other branch, $\mu \geq \alpha > 3b/(2L)$ gives $r = H_*\mu/b > 24L^2 > L$, so affine refinement is legal, and

$$E \leq \frac{2\pi L^2 b N}{H_*} < \frac{\pi b N}{8L} < \frac{b N}{2L} < \frac{\alpha N}{3}.$$

Lemma I.67.1 proves the first bound in (I.68.14).

For the square-root guarantee, put $Y = \sqrt{\alpha r/(6\pi\mu)}$ and $L = \max\{1, \lfloor Y \rfloor\}$. If $L = 1$, then $Y < 2$, and a singleton in A lies in a parent and has length at least $Y/2$. Otherwise $Y^2 \leq r/(6\pi) < r$, so $L \leq r$, and (I.68.16) gives $E \leq \alpha N/3$. The selected cell is parent-contained and has length $L \geq Y/2 = \sqrt{\alpha r/(24\pi\mu)}$.

Finally $\mu \leq 2\tau b$, $r \leq \tau N$, and $r/\tau \geq S/2$. The cube-root minimum in (I.68.14) is at least $(r/(256\tau))^{1/3}$: this is no larger than either entry of that minimum, since $r/(256\tau) \leq N/256 \leq N/8$. It is at least $S^{1/3}/8$. The square-root bound is at least $\sqrt{\alpha r/(48\pi\tau b)} \geq \sqrt{\alpha S/(96\pi b)}$. A singleton supplies the remaining one in the maximum. Selecting a theorem with the largest applicable length guarantee proves (I.68.15). $\square$

[write] **7 October 2026; sources 40–46.** This is source 26's Corollary I.41.3 with its proof (Theorems I.41.1 and I.41.2); source 42's Theorem I.74.1 replaces the cap $(N/8)^{1/3}$ of the first bound by $\sqrt{N/2}$.

Corollary I.68.5 (Retained quadratic and cubic bounds). *For degree at most two, one may replace the length in Theorem F.46.1 by*

$$\max\{1, S^{1/39}/17\}. \quad (\text{I.68.17})$$

For degree at most three, one may replace it by

$$\max\{1, S^{1/4096}/8\}. \quad (\text{I.68.18})$$

Both conclusions retain the density gain $\alpha/3$ over every cyclic modulus.

Proof. For degree two, set $x = S^{1/39}/17$ and $L = \lceil x \rceil$. The singleton branch is unchanged. If $L \geq 2$, then $S > (17L/2)^{39} > L^3$, so the variance branch is valid. In the high-correlation branch,

$$r > \frac{3S}{8L} > \frac{3 \cdot 17^{39}}{2^{42}} L^{38} > 2^{116} L^{38}.$$

The last inequality is $3 \cdot 17^{39} > 2^{158}$. To verify it entirely with integers, Bernoulli's inequality gives $17^{39} = 2^{156}(1 + 1/16)^{39} \geq (55/16)2^{156}$, and $3 \cdot 55/16 > 4$. The quadratic partition and transfer lemmas now apply.

For degree three, set $x = S^{1/4096}/8$ and $L = \lceil x \rceil$. In the nonsingleton case $S > (4L)^{4096} > L^3$. In the high-correlation branch,

$$r > \frac{3S}{8L} > 3 \cdot 2^{8189} L^{4095} > 2^{7834} L^{1878}.$$

Use the sharper cubic budget (I.65.11). The same variance and transfer arguments complete the proof. $\square$

[write] **7 October 2026; sources 40–46.** The quadratic bound (I.68.17) is source 32's Theorem F.37.1; the cubic bound (I.68.18) is new here and answers source 32's degree-three question (item I.19).

The quadratic guarantee also satisfies, for every $S \geq 1$,

$$\max\{1, S^{1/39}/17\} \geq S^{1/64}/8.$$

If the right side is at most one this is immediate. Otherwise $S > 8^{64}$ and $S^{1/39-1/64} > 8^{25/39} > \sqrt{8} > 17/8$; the last comparison is $512 > 289$ after squaring. This illustrates why constants and small scales matter. For any fixed input one may select the largest among the applicable exact-budget, uniform, and small-degree guarantees. The uniform all-degree theorem is not asserted to dominate those earlier bounds pointwise.

I.69 The covariance square and the range inequality

Source 42, delivered Section 2.

We first work in an arbitrary finite abelian group G . For a probability mass function μ on G , let s, t be independent samples with law μ , let a be independent and uniform on G , and put $d = s - t$.

Lemma I.69.1 (Covariance square). *For every real-valued function f on G and every integer h ,*

$$\mathbb{E}_{a,s,t} f(a) f(a + h(s - t)) = \mathbb{E}_z \left(\sum_{s \in G} \mu(s) f(z + hs) \right)^2 \geq 0. \quad (\text{I.69.1})$$

No invertibility assumption on the map $x \mapsto hx$ is needed.

Proof. For each fixed s, t , substitute $a = z + ht$. This is a translation of the uniformly sampled group element, so the left side becomes

$$\mathbb{E}_z \sum_{s,t \in G} \mu(s) \mu(t) f(z + ht) f(z + hs).$$

Independence factors the double sum into the square in (I.69.1). The change of variable only uses translation, never division by h . $\square$

Lemma I.69.2 (Variance and maximum). *If a nonnegative random variable Y satisfies $Y \leq M$ and $\mathbb{E}Y = \delta > 0$, then*

$$\text{Var}(Y) \leq \delta(M - \delta). \quad (\text{I.69.2})$$

Equality holds if and only if $Y \in \{0, M\}$ almost surely.

Proof. The pointwise inequality $Y^2 \leq MY$ gives $\text{Var}(Y) = \mathbb{E}Y^2 - \delta^2 \leq M\delta - \delta^2$. Equality is equivalent to $Y(M - Y) = 0$ almost surely. $\square$

This is the lower-endpoint-zero case of the classical variance bound of Bhatia and Davis [181]. The square identity is likewise a classical autocorrelation-positivity calculation. Neither ingredient is asserted to be new.

Theorem I.69.3 (Admissible difference support). *Let G be finite abelian and $L \geq 2$. Suppose $B \subseteq G$, $|B| = S \geq 2$, and every nonzero element of $B - B$ has additive order at least L . Let $g : G \rightarrow [0, 1]$ have mean $\delta > 0$ and variance $\sigma^2 = \mathbb{E}(g - \delta)^2$. There is a proper L -term progression P such that*

$$\frac{1}{L} \sum_{x \in P} g(x) - \delta \geq \frac{\sigma^2}{\delta} \max \left\{ \frac{S - L}{L(S - 1)}, 0 \right\}. \quad (\text{I.69.3})$$

For $g = \mathbf{1}_A$, this becomes the increment $(1 - \delta) \max\{S - L, 0\}/(L(S - 1))$.

Proof. Choose s, t uniformly from B and a uniformly from G , independently, and define

$$X = \frac{1}{L} \sum_{j=0}^{L-1} g(a + j(s - t)), \quad f = g - \delta.$$

For every fixed s, t , translation invariance gives $\mathbb{E}_a X = \delta$. Expanding the variance yields

$$\text{Var}(X) = \frac{1}{L^2} \sum_{j,k=0}^{L-1} \mathbb{E}_{a,s,t} f(a + j(s - t)) f(a + k(s - t)).$$

Each of the L diagonal terms is σ^2 . For $j \neq k$, translating a and applying Lemma I.69.1 with $h = k - j$ makes the term a nonnegative square. Hence

$$\text{Var}(X) \geq \frac{\sigma^2}{L}. \quad (\text{I.69.4})$$

The event $d = 0$ is precisely $s = t$ and has probability $\eta = 1/S$. On this event, $X = g(a)$ and the conditional variance is σ^2 . On the event $d \neq 0$, every sampled progression is proper by the assumption on $B - B$. Both conditional means are δ , since each fixed step has that translation average. Consequently the law of total variance gives the exact identity

$$\text{Var}(X) = \eta\sigma^2 + (1 - \eta) \text{Var}(X \mid d \neq 0). \quad (\text{I.69.5})$$

Let M be the maximum of the window averages of g over all translates and all nonzero steps in $B - B$. These windows are proper. The maximum exists: distinct points of B supply a nonzero step, and the group is finite. Conditionally on $d \neq 0$, one has $0 \leq X \leq M$. Lemma I.69.2 therefore gives

$$\text{Var}(X \mid d \neq 0) \leq \delta(M - \delta).$$

Combining this with (I.69.4) and (I.69.5), then dividing by $\delta(1 - \eta) > 0$, gives

$$M - \delta \geq \frac{\sigma^2}{\delta} \frac{1/L - \eta}{1 - \eta} = \frac{\sigma^2}{\delta} \frac{S - L}{L(S - 1)}. \quad (\text{I.69.6})$$

Finally, averaging over translates of any nonzero step in $B - B$ gives $M \geq \delta$, which replaces a negative right side by zero. For an indicator, $\sigma^2 = \delta(1 - \delta)$. $\square$

Remark I.69.4 (Nonuniform sampling). The same proof applies when s, t have any common probability distribution supported on an admissible B and at least two points have positive mass. The collision probability becomes

$$\eta = \sum_{x \in B} \mu(x)^2 < 1,$$

and (I.69.6) remains valid. The scalar factor

$$F_L(\eta) := \frac{1/L - \eta}{1 - \eta}, \quad F'_L(\eta) = -\frac{L - 1}{L(1 - \eta)^2} < 0 \quad (\text{I.69.7})$$

is strictly decreasing on $[0, 1)$. Minimizing collisions therefore optimizes this particular variance calculation.

The bounded-function statement records the correct variance dependence. Replacing σ^2/δ by $1 - \delta$ for an arbitrary $[0, 1]$ -valued function would be invalid: a constant function has no positive increment. The case $\sigma^2 = 0$ in (I.69.3) correctly gives zero.

I.70 Exact support optimization and integer scales

Source 42, delivered Section 3.

Return to $G = \mathbb{Z}/N\mathbb{Z}$ and the integers m, q in (F.47.1).

Lemma I.70.1 (Largest admissible support). *The interval of residues*

$$B = \{0, 1, \dots, q - 1\} \subseteq \mathbb{Z}/N\mathbb{Z}$$

has no nonzero difference of additive order below L . Every subset with this property has at most q elements. Moreover, $q - 1$ is the largest integer $D < N$ for which all steps $1, \dots, D$ are proper for length L .

Proof. A nonzero difference of two displayed elements has an integer representative d with $0 < |d| < q$. If its additive order h were below L , then $h \mid N$ and $h \leq m$. The relation $hd = 0$ modulo N says that d is a nonzero multiple of N/h , whence

$$|d| \geq \frac{N}{h} \geq \frac{N}{m} = q,$$

a contradiction.

For the upper bound, let K be the unique subgroup of order m . It has q cosets. Two distinct points in the same coset have nonzero difference in K , whose order divides $m < L$. An admissible support can therefore meet each coset in at most one point, and has size at most q . The displayed interval attains q .

The preceding argument certifies all positive steps smaller than q . The residue represented by q has additive order $m < L$. If $q = N$, the same statement means that N is the zero residue; in either case the largest permissible $D < N$ is $q - 1$. $\square$

Proof of Theorem F.47.1. Apply Theorem I.69.3 with $S = q$ and the support in Lemma I.70.1. The nonzero sampled differences have representatives of absolute value less than q . If $q \leq L$, use the translation-average bound $M_L(A) \geq \delta$. $\square$

For any iid distribution supported on an admissible B , Cauchy–Schwarz gives

$$\eta = \sum_{x \in B} \mu(x)^2 \geq \frac{1}{|B|} \geq \frac{1}{q}.$$

Equality holds for the uniform distribution on the displayed interval. By (I.69.7), this optimizes the support-and-collision form of the argument for every pair (N, L) . This is a statement about the specified method. It does not assert that (F.47.2) is the exact optimum over all sets of every fixed density for each individual pair (N, L) .

I.70.1 The geometric triangular law

A useful factorization-free construction takes $B = \{0, \dots, D\}$ with

$$D \geq 1, \quad (L-1)D < N. \quad (\text{I.70.1})$$

For $0 < |d| \leq D$ and $1 \leq h \leq L-1$, one then has $0 < |hd| < N$. Thus hd is not zero modulo N , and every nonzero supported step is proper. The integer step before reduction modulo N has the triangular law

$$\Pr(s-t=u) = \frac{D+1-|u|}{(D+1)^2}, \quad -D \leq u \leq D. \quad (\text{I.70.2})$$

If distinct representatives give the same residue, their masses are added. Since $D < N$, the zero residue comes only from $s = t$, with mass $1/(D+1)$. Theorem I.69.3 yields

$$H_L(A) \geq \frac{1-\delta}{L} \left(1 - \frac{L-1}{D}\right). \quad (\text{I.70.3})$$

The right side increases with D , so the largest integer permitted by (I.70.1) is

$$D_* := \left\lfloor \frac{N-1}{L-1} \right\rfloor.$$

This gives (F.47.3) after taking the maximum with zero.

The arithmetic theorem never weakens the geometric one. Indeed $m \leq L-1$ gives $N \leq q(L-1)$, and hence

$$\frac{N-1}{L-1} < q, \quad D_* \leq q-1.$$

It can be strictly stronger. For prime N , the arithmetic half-gain holds as soon as $N \geq 2L-1$, well beyond the length range certified by the size-only quadratic condition.

I.70.2 Exact thresholds and rounding

For a requested coefficient $0 < c < 1$ in the increment $c(1-\delta)/L$, the geometric calculation certifies it exactly when

$$D_* \geq \left\lceil \frac{L-1}{1-c} \right\rceil.$$

Equivalently, the sufficient size condition is

$$N \geq (L-1) \left\lceil \frac{L-1}{1-c} \right\rceil + 1. \quad (\text{I.70.4})$$

This is an exact equivalence for this size-only certificate, not a necessary condition for every possible density-increment argument.

Proof of Corollary F.47.2 except sharpness. Strict positivity of (I.70.3) is equivalent to $D_* \geq L$, which holds exactly when $N \geq L(L-1) + 1$. Taking $c = 1/2$ in (I.70.4) gives $N \geq 2(L-1)^2 + 1$. For $N \geq CL^2$ with $C > 1$, the admissible budget satisfies

$$D_* \geq \left\lfloor \frac{N}{L} \right\rfloor \geq \lfloor CL \rfloor \geq CL - 1 \geq C(L-1).$$

For the first inequality, write $D_0 = \lfloor N/L \rfloor$. Then $D_0 \geq 1$ and $(L-1)D_0 \leq N - D_0 < N$, so D_0 is permitted by (I.70.1). The displayed chain implies $1 - (L-1)/D_* \geq 1 - 1/C$, as required. $\square$

For fixed N , the largest length certified by the size-only half-gain condition is

$$1 + \left\lfloor \sqrt{\frac{N-1}{2}} \right\rfloor. \quad (\text{I.70.5})$$

The qualification “size-only” is essential: the arithmetic parameter q can certify a larger length. Formula (I.70.5) describes the certificate $N \geq 2(L-1)^2 + 1$, not the absolute largest length for which all subsets must exhibit that gain.

I.71 Positive definite weights and the classical packing bound

Source 42, delivered Section 4.

The collision optimum extends beyond iid probability differences. We state the normalization explicitly to avoid factors of N . For a probability mass function $w : G \rightarrow [0, 1]$, set

$$W(\chi) = \sum_{d \in G} w(d) \overline{\chi(d)}, \quad \chi \in \widehat{G}.$$

Here $\widehat{G}$ is the character group, and W is an unnormalized Fourier transform. We call w *positive definite* when every $W(\chi)$ is real and nonnegative. The finite Fourier equivalence with the usual positive-semidefinite definition is classical; we only need the displayed condition.

Proposition I.71.1 (Optimal zero-step mass). *Suppose $G = \mathbb{Z}/N\mathbb{Z}$, $N \geq L \geq 2$, and w is a probability mass function such that*

$$W(\chi) \geq 0 \quad (\chi \in \widehat{G}), \quad w(d) = 0 \quad \text{if } d \neq 0 \text{ and } \text{ord}(d) < L.$$

Then

$$w(0) \geq \frac{1}{q}. \quad (\text{I.71.1})$$

Equality is attained by the law of $s - t$ for independent uniform $s, t \in \{0, \dots, q-1\}$.

Proof. Let $K \leq G$ be the subgroup of order m , and let $K^\perp = \{\chi : \chi(k) = 1 \text{ for all } k \in K\}$. Every nonzero $k \in K$ has order below L , so $\sum_{k \in K} w(k) = w(0)$. Fourier inversion and character orthogonality give

$$\begin{aligned} w(0) &= \sum_{k \in K} \frac{1}{N} \sum_{\chi \in \widehat{G}} W(\chi) \chi(k) \\ &= \frac{m}{N} \sum_{\chi \in K^\perp} W(\chi) \geq \frac{m}{N} W(1) = \frac{m}{N} = \frac{1}{q}. \end{aligned}$$

For completeness, inversion follows from $\sum_{\chi \in \widehat{G}} \chi(k-d) = N \mathbf{1}_{k=d}$, and the subgroup sum is m for $\chi \in K^\perp$ and zero otherwise: a nontrivial character on K makes that sum invariant under multiplication by a scalar different from 1. The trivial character has $W(1) = \sum_d w(d) = 1$.

For attainment, let μ be uniform on $\{0, \dots, q-1\}$ and set $w(d) = \sum_x \mu(x) \mu(x-d)$. This law has total mass 1 and

$$W(\chi) = \left| \sum_x \mu(x) \overline{\chi(x)} \right|^2 \geq 0, \quad w(0) = \sum_x \mu(x)^2 = \frac{1}{q}.$$

Lemma I.70.1 verifies the proper nonzero support condition. $\square$

This proposition is a direct specialization of the classical finite-group Turán packing theorem of Kolountzakis and Révész [182, Theorem 2]. In their notation take

$$\Omega = \{0\} \cup \{d : \text{ord}(d) \geq L\}, \quad \Lambda = K.$$

Then $(\Lambda - \Lambda) \cap \Omega = \{0\}$, and their bound reads $\sum_G w \leq (N/m)w(0) = qw(0)$. The short proof above is finite-group Poisson summation. The report claims neither this extremal principle nor the positivity principle as new.

I.71.1 Exactly what is optimized

For a real function f use the normalized transform

$$\widehat{f}(\chi) = \frac{1}{N} \sum_{a \in G} f(a) \overline{\chi(a)}.$$

For an integer h write $(h\chi)(d) = \chi(hd)$. Fourier expansion and the uniform average over a give

$$\mathbb{E}_a \sum_d w(d) f(a) f(a+hd) = \sum_{\chi \in \widehat{G}} |\widehat{f}(\chi)|^2 W(h\chi) \geq 0. \quad (\text{I.71.2})$$

Indeed the a -average pairs inverse characters, and the remaining character factor is $\overline{\chi(hd)}$. This remains a character when multiplication by h is not invertible.

For $f = \mathbf{1}_A - \delta$, the same progression-variance expansion as in Section I.69 gives $\text{Var}(X) \geq \delta(1-\delta)/L$. With $\eta = w(0) < 1$, removing the zero step and applying the range inequality yields

$$H_L(A) \geq (1-\delta) \frac{1/L - \eta}{1-\eta}. \quad (\text{I.71.3})$$

By Proposition I.71.1 and (I.69.7), the largest right side obtained by this exact covariance-positivity, zero-step-removal, and range-bound calculation is attained at $\eta = 1/q$. It is the untruncated arithmetic expression in (F.47.2). The point mass at zero is allowed in Proposition I.71.1, but has $\eta = 1$ and cannot be used in this conditional calculation.

Thus method optimality holds in the entire class of nonnegative probability weights with nonnegative Fourier transform and proper nonzero support. This does not rule out stronger arguments using additional information about A , sharper off-diagonal covariances, or more than its mean and range. For a bounded function g , the same discussion replaces $1 - \delta$ by σ^2/δ .

I.72 Exact extremizers and the optimal uniform constant

Source 42, delivered Section 5.

We now prove all optimality claims about the density increment itself.

Theorem I.72.1 (Prime-quotient equality). *Let p be prime, let $2 \leq L \leq p$, and put*

$$N = p(L - 1), \quad A = \{x \in \mathbb{Z}/N\mathbb{Z} : x \equiv 0 \pmod{p}\}.$$

Then $\delta = 1/p$, $m(N, L) = L - 1$, $q(N, L) = p$, and

$$M_L(A) = \frac{1}{L}, \quad H_L(A) = \frac{p - L}{Lp} = \frac{1 - \delta}{L} \frac{p - L}{p - 1}. \quad (\text{I.72.1})$$

Hence the arithmetic theorem is attained exactly throughout this family.

Proof. The set consists of $L - 1$ multiples of p modulo N , so it has density $1/p$. The integer $L - 1$ divides N and is the largest possible integer below L , giving $m = L - 1$.

If $p \mid d$, the additive order of d modulo N divides $L - 1$. Such a step cannot be proper for length L . For every proper step, therefore, d is nonzero modulo p . Since p is prime and $L \leq p$, the first L residues

$$a, a + d, \dots, a + (L - 1)d \pmod{p}$$

are distinct. At most one is zero, so every proper progression meets A in at most one point. Conversely, $P(0, 1; L)$ is proper because $N \geq L$, and contains precisely one point of A . Thus $M_L(A) = 1/L$ and (I.72.1) follows. $\square$

For $L < p$, equality can also be seen inside the proof. Use $B = \{0, \dots, p - 1\}$. For every nonzero difference of indices h with $|h| < L \leq p$, multiplication by h permutes residues modulo p . The local average of $\mathbf{1}_A - 1/p$ in the covariance square is therefore zero. Every nonzero sampled window has average 0 or $1/L$, giving equality in the range bound. The endpoint $L = p$ has zero increment; all proper progressions then meet A exactly once.

Corollary I.72.2 (Sharp positivity threshold and failure at coefficient one). *For infinitely many L , at $N = L(L - 1)$ there is a nonempty proper subset with no positive density increment on any proper L -term progression. No condition $N \geq cL^\gamma$ with fixed $c > 0$ and $\gamma < 2$ can guarantee such a positive increment uniformly. Further, the condition $N \geq L^2$ does not guarantee $H_L(A) \geq c(1 - \delta)/L$ for any absolute $c > 0$.*

Proof. Take $L = p$ in Theorem I.72.1. Then $N = p(p - 1)$ and $H_L(A) = 0$. As primes tend to infinity, $p(p - 1)/(p^\gamma) \rightarrow \infty$ for every $\gamma < 2$, proving the second assertion.

For the last assertion take primes $p \geq 5$ and $L = p - 2$. Then

$$N = p(p - 3) = L^2 + L - 2 > L^2, \quad \frac{LH_L(A)}{1 - \delta} = \frac{2}{p - 1} \rightarrow 0.$$

An absolute positive coefficient is therefore impossible even with the strict inequality $N > L^2$. $\square$

Theorem I.72.3 (Optimal quadratic size–gain tradeoff). *Fix a real number $C > 1$. The largest real c for which*

$$N \geq CL^2, \quad N \geq L \geq 2, \quad 0 < |A| < N \quad \implies \quad H_L(A) \geq c \frac{1 - |A|/N}{L}$$

holds for every such N, L, A is exactly $c = 1 - 1/C$.

Proof. Corollary F.47.2 proves that $c = 1 - 1/C$ works. To rule out every larger value while satisfying the exact nonasymptotic size hypothesis, let p tend to infinity through primes and set

$$L = \left\lfloor \frac{p}{C} \right\rfloor - 2, \quad N = p(L - 1).$$

Eventually $L \geq 2$, and $L \leq p$. Moreover $p \geq C(L + 2)$, so

$$N \geq C(L + 2)(L - 1) = C(L^2 + L - 2) \geq CL^2.$$

Use the sets of Theorem I.72.1. Since $L/p \rightarrow 1/C$,

$$\frac{LH_L(A)}{1 - \delta} = \frac{p - L}{p - 1} \rightarrow 1 - \frac{1}{C}.$$

Any proposed coefficient larger than this limit fails for all sufficiently large members of the sequence. Only the infinitude of primes is used, not a prime-gap estimate. $\square$

The theorem optimizes a bound over *all* nontrivial densities. Its extremizers satisfy $\delta = 1/p \rightarrow 0$. Nothing in this argument identifies the best coefficient when δ is held fixed or bounded below. In particular, the phrase “sharp quadratic scale” refers to the explicit arithmetic equalities and the stated uniform size–gain tradeoff, not to a complete extremal classification.

I.73 A density sensitive gain at the quadratic endpoint

Source 42, delivered Section 6.

The uniform coefficient in Theorem I.72.3 tends to zero as C decreases to 1. That conclusion is driven by sparse examples. A different elementary argument gives a positive coefficient at the endpoint when density is kept away from zero and one.

Theorem I.73.1 (Sliding-window endpoint bound). *Let $L \geq 2$, $N > L(L - 1)$, and $0 < \delta = |A|/N < 1$. Define the proportion of short-order directions by*

$$\beta = \frac{1}{N} \sum_{\substack{h|N \\ h < L}} \varphi(h),$$

where φ is Euler's totient function. Then

$$H_L(A) \geq \frac{\delta(1-\delta)}{L(1+\beta)+\delta} > \frac{\delta(1-\delta)}{2L}. \quad (\text{I.73.1})$$

In particular, if $0 < \varepsilon \leq 1/2$ and $\varepsilon \leq \delta \leq 1-\varepsilon$, the density gain exceeds $\varepsilon(1-\varepsilon)/(2L)$. No optimality is claimed for this coefficient.

Proof. Let $\mathcal{P} = \{d \in G : \text{ord}(d) \geq L\}$ be the proper directions. Set

$$T_d(a) = \sum_{j=0}^{L-1} \mathbf{1}_A(a + jd), \quad M = \max_{a, d \in \mathcal{P}} T_d(a), \quad \eta = M - \delta L = LH_L(A).$$

Averaging over a for any proper direction gives $\eta \geq 0$. For each $d \in \mathcal{P}$, the deficits $D_d(a) = M - T_d(a)$ are nonnegative and have mean η . The sliding identity is

$$T_d(a + d) - T_d(a) = \mathbf{1}_A(a + Ld) - \mathbf{1}_A(a).$$

Consequently

$$|\mathbf{1}_A(a + Ld) - \mathbf{1}_A(a)| = |D_d(a + d) - D_d(a)| \leq D_d(a + d) + D_d(a).$$

Define $F(x) = |A \triangle (A + Lx)|/(2N)$ for every $x \in G$. After averaging over a , the preceding inequality becomes

$$F(d) \leq \eta \quad (d \in \mathcal{P}). \quad (\text{I.73.2})$$

A cyclic group has $\varphi(h)$ elements of order h when $h \mid N$ and none otherwise. Indeed such elements correspond to generators of its unique subgroup of order h . Thus

$$|G \setminus \mathcal{P}| = \sum_{\substack{h \mid N \\ h < L}} \varphi(h) \leq \sum_{h=1}^{L-1} h = \frac{L(L-1)}{2} < \frac{N}{2}.$$

Hence $|\mathcal{P}| > N/2$ and $\mathcal{P} - \mathcal{P} = G$: for every x , the sets $\mathcal{P}$ and $x + \mathcal{P}$ intersect. Translation invariance and the triangle inequality for symmetric difference give $F(x+y) \leq F(x) + F(y)$ and $F(-x) = F(x)$. Writing $x = d_1 - d_2$ with $d_1, d_2 \in \mathcal{P}$, (I.73.2) therefore gives $F(x) \leq 2\eta$ for every $x \in G$. Averaging the better bound on $\mathcal{P}$ and this bound on its complement yields

$$\mathbb{E}_x F(x) \leq (1-\beta)\eta + 2\beta\eta = (1+\beta)\eta. \quad (\text{I.73.3})$$

Put $g = \gcd(N, L)$ and $K = LG = gG$, so $|K| = N/g \geq N/L > L-1$. The integer $|K|$ is therefore at least L . Let $p_C = |A \cap C|/|K|$ for each coset C of K , and average equally over these cosets. Then $\mathbb{E}_C p_C = \delta$. The variable Lx , for uniform $x \in G$, is uniform on K ; averaging the symmetric difference within each coset gives the exact identity

$$\mathbb{E}_x F(x) = \mathbb{E}_C p_C(1 - p_C). \quad (\text{I.73.4})$$

For example, the averaged normalized number of ordered pairs with the first point in $A \cap C$ and the second in $C \setminus A$ is $p_C(1 - p_C)$; the symmetric difference counts both orientations and its denominator contains the factor 2.

The step g has order $|K| \geq L$, so it is proper. Averaging its L -window counts over starting points in any coset C gives $Lp_C \leq M$. It follows that

$$p_C \leq \delta + \frac{\eta}{L}, \quad \mathbb{E}_C p_C(1 - p_C) = \delta - \mathbb{E}_C p_C^2 \geq \delta(1 - \delta) - \frac{\delta\eta}{L}.$$

Combining this with (I.73.3)–(I.73.4) gives

$$\eta \left(1 + \beta + \frac{\delta}{L}\right) \geq \delta(1 - \delta).$$

Since $\beta < 1/2$ and $\delta/L < 1/2$, the factor in parentheses is strictly below 2. Division by L proves (I.73.1); the stated density-range corollary follows from $\delta(1 - \delta) \geq \varepsilon(1 - \varepsilon)$. $\square$

The proof only needs $|\mathcal{P}| > N/2$ and $N/\gcd(N, L) \geq L$; the displayed size condition is a convenient sufficient condition for both. Unlike the positive-definite calculation, this argument uses the maximum of all windows to force approximate invariance under L -multiple shifts. Its factor δ makes it compatible with the vanishing-density equality examples in Section I.72.

Proposition I.73.2 (Dense examples at the strict boundary). *Let p be prime, let $L = p$ and $N = p(p - 1)$, and let A be the inverse image of any k -element subset of $\mathbb{Z}/p\mathbb{Z}$, $0 < k < p$. Then every proper L -term progression has density exactly $\delta = k/p$.*

Proof. A step divisible by p has order at most $p - 1$ and is not proper. Every proper step is nonzero modulo p , so its p projected points visit every residue modulo p exactly once. Exactly k points belong to the prescribed inverse image. $\square$

Choosing $k = \lfloor \delta_0 p \rfloor$ for any fixed $0 < \delta_0 < 1$ gives densities tending to δ_0 and no positive gain at $N = L(L - 1)$. Thus the strict universal positivity threshold is sharp even for densities confined to a fixed interval away from zero and one. This does not prove sharpness of the quantitative constant in Theorem I.73.1.

The cyclic size condition cannot simply be extended to all finite abelian groups by replacing N with group cardinality. For a fixed prime $p \geq 3$, take $G = (\mathbb{Z}/(p - 1)\mathbb{Z})^s \times \mathbb{Z}/p\mathbb{Z}$ and let A be the inverse image of a nonempty proper subset of the last factor. A direction with zero last coordinate has order dividing $p - 1$; every proper p -term progression therefore projects bijectively onto the last factor and has exactly the global density. The group size tends to infinity with s , without a positive increment. The admissible-support hypothesis of Theorem I.69.3 remains essential.

I.74 A sparse affine cover consequence

Source 42, delivered Section 7.

We give a self-contained application in which the new scale changes a concrete part of a density-increment conclusion. Write $e(t) = \exp(2\pi it)$, and $e_N(y) = e(\tilde{y}/N)$ for any integer representative of $y \in \mathbb{Z}/N\mathbb{Z}$.

Theorem I.74.1 (Affine exact-mass bound). *Let $N \geq 2$ and let $A \subseteq \mathbb{Z}/N\mathbb{Z}$ have density δ . Put $f = \mathbf{1}_A - \delta$. Let $M \geq 1$ and let $P_1, \dots, P_M$ be nonempty pairwise disjoint proper modular progressions*

with $A \subseteq U := \bigcup_i P_i$. Let $\phi_i(x) = a_i x + b_i$ be affine maps with coefficients in $\mathbb{Z}/N\mathbb{Z}$, and suppose $\alpha > 0$ and

$$\sum_{i=1}^M \left| \sum_{x \in P_i} f(x) e_N(-\phi_i(x)) \right| \geq \alpha N. \quad (\text{I.74.1})$$

Set

$$r = \min_i |P_i|, \quad \mu = \frac{1}{N} \sum_{x \in U} |f(x)|, \quad b = 1 - \delta.$$

Then some proper modular progression Q satisfies

$$|Q| \geq \max \left\{ 1, \min \left\{ \left(\frac{rb}{128\mu} \right)^{1/3}, \sqrt{\frac{N}{2}} \right\} \right\}, \quad \text{dens}_Q(A) \geq \delta + \frac{\alpha}{3}. \quad (\text{I.74.2})$$

No comparability assumption on parent lengths is required. The overall conclusion does not assert $Q \subseteq U$.

[write] **7 October 2026; sources 40–46.** Against source 26 (checked at the write). Source 26’s Theorem I.41.1 gives length $\min\{(r(1-\delta)/(128\mu))^{1/3}, (N/8)^{1/3}\}$; here the first term is the same and only the modulus cap $(N/8)^{1/3}$, which comes from the sufficient condition $N \geq L^3$ of the variance step, becomes $\sqrt{N}/2$. Since $\sqrt{N}/2/(N/8)^{1/3} = \sqrt{2} N^{1/6} > 1$, the new bound is never smaller, and it is larger exactly when the cap binds, that is when $r > 16\mu N/(1-\delta)$ (sparse sets; hence the section title). So “26’s affine cap $(N/8)^{1/3}$ becomes $\sqrt{N}/2$ ” is the precise reading of the placement’s “square-root cap”; the cube-root dependence on r is not changed. This answers the first of source 26’s questions (item I.19).

The corresponding exact-mass statement of Report275 [195], retained in Report277, has the modulus cap $(N/8)^{1/3}$ instead of $\sqrt{N}/2$. The first term in (I.74.2) and its gain are unchanged. Thus this consequence improves the independent variance cap, without changing a uniform comparable-cover exponent or any polynomial phase exponent.

I.74.1 Mass identities and affine refinement

Let $\tau = |U|/N$. The correlation hypothesis implies $0 < \delta < 1$; otherwise $f = 0$. Counting A and $U \setminus A$ gives

$$0 < \alpha \leq \mu = \delta(1 + \tau - 2\delta) \leq 2\delta b, \quad \sum_{x \in U} f(x) = \delta N(1 - \tau) \geq 0. \quad (\text{I.74.3})$$

The first inequality also uses the triangle inequality and $\tau \leq 1$. In particular, all quantities in (I.74.2) are well-defined and positive.

Lemma I.74.2 (Affine cells). *For an integer $1 \leq L \leq r$, the parents admit a partition into proper progression cells C with $L \leq |C| \leq 2L - 1$. Each cell has a unit complex number z_C such that*

$$E := \sum_C \sum_{x \in C} |f(x)| |e_N(-\phi_{i(C)}(x)) - z_C| \leq \frac{2\pi L^2 \mu N}{r}. \quad (\text{I.74.4})$$

If $E \leq \alpha N/3$, one of the cells has density at least $\delta + \alpha/3$.

Proof. Write a parent as $u + jd$, $0 \leq j < n$, with injective parameterization. Its affine phase is a constant multiple of $e(-j\theta)$ for a real representative θ of $a_i d/N$. Put $R = \lfloor n/L \rfloor \geq 1$. Among the $R + 1$ points $0, \theta, \dots, R\theta$ on the circle, some circular gap has length at most $1/(R + 1)$. Subtract

the corresponding two indices and change sign if necessary. This produces integers $1 \leq p \leq R$ and v with

$$p\theta = v + \varepsilon, \quad |\varepsilon| \leq \frac{1}{R+1}. \quad (\text{I.74.5})$$

Coincident points give $\varepsilon = 0$; the circular gap crossing zero is included. This is the elementary circle pigeonhole estimate.

Partition the parameter interval into increasing residue chains modulo p . Each chain has at least $\lfloor n/p \rfloor \geq L$ elements. A chain of length $T = kL + s$, $0 \leq s < L$, splits into $k-1$ chunks of length L and one of length $L+s$. Each chunk maps to a proper modular progression because the full parent parameterization is injective.

On a chunk of length ℓ , write its phase as a constant multiple of $e(-j\varepsilon)$, $0 \leq j < \ell$, using (I.74.5). Choose z_C to be the phase at the real midpoint $(\ell-1)/2$ of this *reduced* phase. The elementary chord-arc bound $|e(t) - e(u)| \leq 2\pi|t - u|$ gives

$$|e_N(-\phi_i(x)) - z_C| \leq 2\pi|\varepsilon| \left| j - \frac{\ell-1}{2} \right| \leq \frac{2\pi(L-1)}{R+1} < \frac{2\pi L^2}{n}$$

when $L \geq 2$. For $L = 1$, the cells are singletons and the error is zero. The reduction of the integer v before evaluating the midpoint is important: a half-integral midpoint must not inadvertently retain an integer phase increment. Summing the pointwise estimate over parents, using $n \geq r$, proves (I.74.4).

Put $F_C = \sum_{x \in C} f(x)$. Replacing the phases by z_C and using the triangle inequality gives

$$\sum_C |F_C| \geq \alpha N - E.$$

Also $\sum_C F_C = \sum_U f \geq 0$ by (I.74.3). Thus

$$\sum_{F_C > 0} F_C = \frac{1}{2} \left(\sum_C |F_C| + \sum_C F_C \right) \geq \frac{\alpha N - E}{2}.$$

If $E \leq \alpha N/3$, the positive cells carry balanced mass at least $\alpha N/3$ and occupy at most N points. At least one consequently satisfies $F_C/|C| \geq \alpha/3$. $\square$

I.74.2 The two branches with exact ceiling control

Proof of Theorem I.74.1. Set

$$H = \frac{rb}{\mu}, \quad x = \min\{(H/128)^{1/3}, \sqrt{N/2}\}, \quad L = \lceil x \rceil.$$

If $L = 1$, choose a singleton in A . This has the requested length, and $\alpha/3 \leq 2\delta b/3 < b$ by (I.74.3), so its density 1 exceeds $\delta + \alpha/3$.

Suppose $L \geq 2$. Then $x > 1$, $L < 2x$, and

$$H \geq 128x^3 > 16L^3. \quad (\text{I.74.6})$$

Independently, the definition of L gives

$$L-1 < x \leq \sqrt{N/2}, \quad 2(L-1)^2 < N.$$

Both N and $2(L-1)^2$ are integers, hence

$$N \geq 2(L-1)^2 + 1. \quad (\text{I.74.7})$$

In particular $N \geq L$, and the quadratic half-gain corollary applies.

If $\alpha \leq 3b/(2L)$, Corollary F.47.2 yields a proper progression of length L with density at least $\delta + b/(2L) \geq \delta + \alpha/3$. Since $L \geq x$ and $L \geq 1$, this proves (I.74.2) in the low-correlation branch.

Otherwise $\alpha > 3b/(2L)$. By (I.74.3) and (I.74.6),

$$r = \frac{H\mu}{b} > \frac{16L^3\mu}{b} > 24L^2 > L.$$

Thus Lemma I.74.2 is applicable. Its error satisfies

$$E \leq \frac{2\pi L^2\mu N}{r} = \frac{2\pi L^2bN}{H} < \frac{\pi bN}{8L} < \frac{bN}{2L} < \frac{\alpha N}{3},$$

where $\pi < 4$ was used. The lemma supplies a proper cell of length at least L with the required density. This proves the high-correlation branch and the theorem. $\square$

The high-correlation cell is contained in a parent. The low-correlation progression comes from the variance theorem and need not lie in U . For that reason the theorem as a whole carries no containment assertion. Nor does properness imply that the progression is nonwrapping when represented in an ordinary integer interval.

I.75 Source 67's Section 2: unequal Fourier budgets at Lemmas 5.2 and 5.5 [write]

Added in the eleventh write, 7 October 2026.

The next section is source 67's Section 2. *Logically it continues source 37's chapter* on the missing-interval step of Gowers's Section 5 (Section I.49); it is printed here, at the end of Part I, because no number printed earlier may move. Source 67's introduction is Section F.72; its other sections are in Parts II, III and V.

What is proved. Source 37 found the least mass of the open arc I_a for a probability measure on the circle whose first R Fourier coefficients are bounded by one common budget $e \leq \rho$, namely $(\rho - e)/(1 + \rho)$ in the window $\pi/(R+1) < a < 2\pi/(R+1)$ (Theorem I.50.1). With unequal budgets $|\hat{\mu}(j)| \leq \varepsilon_j$ the least mass is $V_{R,a}(\boldsymbol{\varepsilon}) = -1/((1 - \cos a) \mathbf{e}^T L(\boldsymbol{\varepsilon})^{-1} \mathbf{e})$ under finitely checkable strict hypotheses on the localizing matrix L (Theorem I.76.2), which hold on an open neighbourhood of every interior point $e\mathbf{e}$, $0 < e < \rho$, of the equal-budget branch (Theorem I.76.4). Source 37's certificate stays optimal exactly on an explicit polytope of local dimension $\lfloor R/2 \rfloor + 1$ (Theorem I.76.5); at every strict-certificate point the Hessian of the value has rank $\lfloor (R-1)/2 \rfloor$ (Theorem I.76.6), so for $R \geq 3$ unequal budgets usually need a moving certificate, with an exact rational example at $R = 3$, $a = \pi/3$ (Corollary I.76.7); on aligned grids the elementary approximation improves to order N^{-2} (Theorem I.76.8). The supported Toeplitz moment criterion is Yang and Xie's, as for source 37, and is credited.

Relation to the formal statements (checked at HEAD fb9602e55). `lemma_5_2` (Section05.lean, line 141) has the exact companion `lemma_5_2_holds` (Proofs05MissingInterval.lean, line 220); source 67's formula concerns the extremal problem behind it, not a different statement, and is not formalized.

I.76 Exact Fourier-gap profiles with unequal budgets

Source 67, delivered Section 2.

The missing-interval step in Gowers' Lemma 5.2 converts a gap in the support of a measure into a large Fourier coefficient. Source 37 of the repository obtained the exact answer when the first several coefficients have a common bound. Its Question 4 asks for unequal bounds and for the region on which the old certificate remains optimal. We give a matrix formula on an open neighbourhood of every interior equal-budget point, a finite certificate valid beyond those neighbourhoods, and an exact description of the old certificate's entire optimality region. These statements concern genuine optima over measures, rather than the output of a discretized optimization.

I.76.1 The extremal problem and its moment matrices

Write $\mathbb{T} = \mathbb{R}/(2\pi\mathbb{Z})$, and, for $0 < a < \pi$, put

$$I_a = \{\theta : |\theta|_o < a\}, \quad K_a = \mathbb{T} \setminus I_a, \quad c = \cos a.$$

The excluded arc is open. Its endpoints are allowed in K_a . For an integer $R \geq 1$ and a vector $\varepsilon = (\varepsilon_1, \dots, \varepsilon_R) \geq 0$, use $\hat{\mu}(j) = \int e^{-ij\theta} d\mu(\theta)$ and define

$$V_{R,a}(\varepsilon) = \inf \{\mu(I_a) : \mu \text{ a probability measure on } \mathbb{T}, |\hat{\mu}(j)| \leq \varepsilon_j \ (1 \leq j \leq R)\}. \quad (\text{I.76.1})$$

Reflection followed by averaging preserves arc mass and does not increase any Fourier modulus, so a symmetric minimizer may be sought. The uniqueness conclusions below nevertheless apply to all probability measures. Minimizers exist: probability measures on the compact circle form a weakly compact space, the moment constraints are closed, and the mass of the open set I_a is lower semicontinuous by the Portmanteau theorem.

Set $h_0 = 1$ and $h_j = h_{-j} = -\varepsilon_j$ for $1 \leq j \leq R$. The following real symmetric matrices depend affinely on the budgets:

$$M(\varepsilon) = (h_{j-k})_{0 \leq j,k \leq R}, \quad (\text{I.76.2})$$

$$L(\varepsilon) = \left(ch_{j-k} - \frac{h_{j-k+1} + h_{j-k-1}}{2} \right)_{0 \leq j,k < R}. \quad (\text{I.76.3})$$

Let $\mathbf{e}$ denote the all-ones vector in $\mathbb{R}^R$ and $Z = \mathbf{e}^\perp$. We say that L is positive on Z when $v^T L v > 0$ for every nonzero $v \in Z$; for $R = 1$ this condition is vacuous.

We use a standard supported moment theorem, in the normalization needed here. The matrix criterion is classical; see Yang and Xie [152].

Lemma I.76.1 (Moment criterion for the complementary arc). *Given $g_{-R}, \dots, g_R$ with $g_{-j} = \overline{g_j}$ and $g_0 > 0$, there is a positive measure of mass g_0 supported on K_a with moments g_j if and only if*

$$(g_{j-k})_{0 \leq j,k \leq R} \succeq 0, \quad \left(cg_{j-k} - \frac{g_{j-k+1} + g_{j-k-1}}{2} \right)_{0 \leq j,k < R} \succeq 0.$$

Proof. Necessity follows by integrating $|A(e^{i\theta})|^2$ and $(c - \cos \theta)|B(e^{i\theta})|^2$. For sufficiency, the interval factorization for nonnegative real polynomials gives the following one-arc representation: every real trigonometric polynomial H of degree at most R which is nonnegative on K_a has the form

$$H(\theta) = |A(e^{i\theta})|^2 + (c - \cos \theta)|B(e^{i\theta})|^2, \quad \deg A \leq R, \quad \deg B \leq R - 1.$$

Indeed, $t = \cot(\theta/2)$ maps K_a to $[-b, b]$, $b = \cot(a/2)$. Multiply H by $(1 + t^2)^R$ and apply the real interval factorization $p(t) = U(t)^2 + (b^2 - t^2)W(t)^2$, with degrees at most R and $R - 1$. Dividing back and using $(b^2 - t^2)/(1 + t^2) = (c - \cos \theta)/(1 - c)$ proves the representation. This is the usual Markov–Lukács factorization; the one-arc criterion can equivalently be imported directly from [152].

The prescribed moments define a real linear functional on degree- R trigonometric polynomials. The two matrix inequalities and the representation make it nonnegative on every polynomial nonnegative on K_a . If the normalized moment vector were outside the convex hull of $\theta \mapsto (\cos \theta, -\sin \theta, \dots, \cos R\theta, -\sin R\theta)$ for $\theta \in K_a$, strict separation would contradict that positivity. The convex hull is compact by finite-dimensional Carathéodory, so this argument supplies a positive representing measure. $\square$

I.76.2 A finite matrix certificate for the exact answer

Theorem I.76.2 (Exact unequal-budget formula). *Suppose the following conditions hold:*

(i) $L = L(\varepsilon)$ is positive on Z , is invertible, and $\beta = \mathbf{e}^T L^{-1} \mathbf{e} < 0$.

(ii) For

$$s = -\frac{1}{(1 - c)\beta}, \quad w = \frac{L^{-1} \mathbf{e}}{\beta}, \quad P_w(z) = \sum_{j=0}^{R-1} w_j z^j, \quad (\text{I.76.4})$$

one has $0 < s < 1$ and $M(\varepsilon) - sJ_{R+1} \succ 0$.

(iii) In the expansion

$$F_w(\theta) = \frac{(\cos \theta - c)|P_w(e^{i\theta})|^2}{1 - c} = f_0 + 2 \sum_{j=1}^R f_j \cos(j\theta), \quad (\text{I.76.5})$$

every f_j with $1 \leq j \leq R$ is strictly positive.

Then

$$V_{R,a}(\varepsilon) = -\frac{1}{(1 - \cos a) \mathbf{e}^T L(\varepsilon)^{-1} \mathbf{e}}. \quad (\text{I.76.6})$$

There is a unique minimizing measure. It has an atom of mass s at 0, exactly $R + 1$ further atoms in K_a , and moments $\hat{\mu}(j) = -\varepsilon_j$ for $1 \leq j \leq R$. The other atoms are $-a$, a , and the arguments of the $R - 1$ zeros of P_w ; those zeros are simple, are on the unit circle, and are different from the endpoints. All $R + 2$ masses are strictly positive.

Proof. Put $\kappa = \beta^{-1} = -(1 - c)s$. Then $\mathbf{e}^T w = 1$ and $Lw = \kappa \mathbf{e}$. Thus

$$K = L + s(1 - c)J_R = L - \kappa J_R$$

annihilates w . Every x decomposes uniquely as $x = u + (\mathbf{e}^T x)w$, with $u \in Z$, and $x^T K x = u^T L u$. Consequently $K \succeq 0$ has kernel $\mathbb{R}w$ and rank $R - 1$.

Subtracting an atom of mass s at zero from the prescribed moments changes M to $M - sJ_{R+1}$ and L to $L + s(1 - c)J_R$. The former is positive definite and the latter positive semidefinite. Lemma I.76.1 therefore gives a positive measure σ of mass $1 - s$ on K_a with moments $h_j - s$. Adding $s\delta_0$ gives a feasible measure μ .

Since $P_w(1) = 1$, $F_w(0) = 1$. Its positive nonconstant Fourier coefficients give $F_w(\theta) \leq 1$ everywhere, with equality only at 0. Also $F_w \leq 0$ on K_a , so $F_w \leq \mathbf{1}_{I_a}$. For any feasible measure ν ,

$$\nu(I_a) \geq f_0 + 2 \sum_{j=1}^R f_j \Re \hat{\nu}(j) \geq f_0 - 2 \sum_{j=1}^R f_j \varepsilon_j = -\frac{w^T L w}{1-c} = s.$$

For the constructed measure, the localizing quadratic form at w vanishes:

$$\int_{K_a} (c - \cos \theta) |P_w(e^{i\theta})|^2 d\sigma(\theta) = 0.$$

Hence its support consists only of endpoints and unit-circle zeros of P_w . The coefficient $f_R > 0$ implies $w_0 w_{R-1} > 0$, so P_w has degree exactly $R-1$. At most $R+1$ such contact points are possible. On the other hand, the positive definite matrix $M - sJ_{R+1}$ has rank $R+1$, and a measure on fewer points cannot have this rank. There are therefore exactly $R+1$ contact points, each with positive mass. This also proves all the root assertions without a numerical root assumption.

Equality for any feasible ν forces every moment to be $-\varepsilon_j$, because each $f_j > 0$, and forces support on the same $R+2$ contact points, including 0. The moments of orders $-R, \dots, R$ determine weights on this support: the corresponding Vandermonde matrix has full column rank since $R+2 \leq 2R+1$. Thus $\nu = \mu$. $\square$

All tests in this theorem involve matrices of order at most $R+1$ and the coefficients of one polynomial. They are sufficient conditions, not an assertion that one rational branch describes every budget vector. At a boundary, a mass, coefficient, or eigenvalue can vanish and the contact pattern can change.

I.76.3 An open region around every interior equal-budget point

For the rest of this subsection suppose

$$\frac{\pi}{R+1} < a < \frac{2\pi}{R+1}, \quad 0 < a < \pi. \quad (\text{I.76.7})$$

Write $m = R+1$ and

$$S = m - \cot(a/2) \tan(ma/2), \quad d = S^{-1}, \quad \rho = \frac{d}{1-d}. \quad (\text{I.76.8})$$

The equal-budget result of source 37 is $V_{R,a}(e\mathbf{e}) = (\rho - e)/(1 + \rho)$ for $0 \leq e \leq \rho$. We recall enough of its certificate to verify the hypotheses of our extension rather than treating their persistence as an assumption.

Lemma I.76.3 (The equal-budget certificate). *Let T_R have $1/2$ in the two adjacent off-diagonals and zero elsewhere. The vector*

$$w_j^0 = \frac{1 - \cos((j - (R-1)/2)a) / \cos(ma/2)}{S}, \quad 0 \leq j < R, \quad (\text{I.76.9})$$

has positive entries, has sum one, and satisfies

$$(T_R - cI)w^0 = (1-c)d\mathbf{e}.$$

The polynomial F_{w^0} has constant coefficient d and all its other Fourier coefficients are positive. Also

$$v^T T_R v \leq \cos(2\pi/m) \|v\|_2^2 \quad (v \in Z). \quad (\text{I.76.10})$$

Proof. In the strict window $\cos(ma/2) < 0$. The cosine in each numerator is larger than that denominator, which proves positivity. Summing a finite cosine progression gives the normalization, and the cosine recurrence, extended by $w_{-1}^0 = w_R^0 = 0$, gives the displayed linear equation. Let $A_k = \sum_j w_j^0 w_{j+k}^0$, with zero extension. The coefficient q_k of $e^{ik\theta}$ in $(\cos \theta - c)|P_{w^0}|^2$ is

$$q_0 = (1 - c)d, \quad q_k = (1 - c)d \sum_{j=0}^{R-k-1} w_j^0 + \frac{1}{2}w_{R-k}^0 w_{R-1}^0 > 0 \quad (1 \leq k \leq R).$$

Division by $1 - c$ proves the coefficient assertion. For (I.76.10), append a zero to v and use the Fourier eigenvalues $\cos(2\pi\ell/m)$ of the cyclic nearest-neighbour matrix. The appended vector still has sum zero, so its constant Fourier mode is absent. $\square$

Theorem I.76.4 (Unequal-budget neighbourhoods). *For every R , a satisfying (I.76.7) and every $e \in (0, \rho)$, there is an open neighbourhood $\mathcal{U}$ of $e\mathbf{e}$ in $\mathbb{R}^R$ on which all the hypotheses of Theorem I.76.2 hold. In particular the exact optimum is a real analytic rational function of the budgets on $\mathcal{U}$.*

Proof. At equal budgets,

$$L = (1 + e)(cI - T_R) + e(1 - c)J_R.$$

On Z this is positive definite by (I.76.10) and $c > \cos(2\pi/m)$. The vector w^0 satisfies

$$Lw^0 = (1 - c)(e - (1 + e)d)\mathbf{e} = -(1 - c)\frac{\rho - e}{1 + \rho}\mathbf{e}.$$

This scalar is negative, so L is invertible and $\beta < 0$. With $s = (\rho - e)/(1 + \rho)$,

$$M - sJ_m = (1 + e)I_m - (e + s)J_m.$$

Its constant-mode eigenvalue is $(1 + e)(1 - md) > 0$ since $d < 1/m$, and its other eigenvalues are $1 + e$. The coefficient condition follows from Lemma I.76.3. All these inequalities are strict and persist in a sufficiently small neighbourhood; budgets there can also be kept positive. $\square$

I.76.4 The exact region for the old certificate

Fix the equal-budget certificate $F_0 = F_{w^0}$, and let

$$Z_0 = \{0, -a, a\} \cup \{\arg z : P_{w^0}(z) = 0\}.$$

It has $R + 2$ points. Write $F_0 = d + 2 \sum f_j^0 \cos(j\theta)$.

Theorem I.76.5 (Complete fixed-certificate optimality region). *For an arbitrary nonnegative budget vector, the lower bound*

$$V_{R,a}(\epsilon) \geq d - 2 \sum_{j=1}^R f_j^0 \epsilon_j \tag{I.76.11}$$

is an equality if and only if the finite system

$$p_z \geq 0, \quad \sum_{z \in Z_0} p_z = 1, \quad \sum_{z \in Z_0} p_z e^{-ijz} = -\epsilon_j \quad (1 \leq j \leq R) \tag{I.76.12}$$

is feasible. Its solution, when it exists, is unique. Thus the exact optimality region is a computable polytope intersected with the nonnegative budget orthant. Near an interior equal-budget point its dimension is $\lfloor R/2 \rfloor + 1$ and its codimension is $\lfloor (R - 1)/2 \rfloor$.

Proof. The certificate proves (I.76.11). If equality holds, each positive coefficient forces the corresponding moment to be $-\varepsilon_j$ and the measure to be supported on the contact set Z_0 . Conversely a measure in (I.76.12) has exactly the certificate's integral as its arc mass, because $F_0 = \mathbf{1}_{I_a}$ on Z_0 . Uniqueness follows from the same Vandermonde argument as before. Reflection consequently fixes the measure.

There are $\lfloor R/2 \rfloor + 2$ reflection orbits in Z_0 : zero, endpoint pairs and interior conjugate pairs, with an additional point at π when R is even. This count follows because P_{w^0} is palindromic, has simple unit-circle roots, and satisfies $P_{w^0}(1) = 1$: odd degree forces a root at -1 , whereas an even-degree palindromic polynomial cannot have a simple root at -1 . Thus the symmetric probability weights have dimension $\lfloor R/2 \rfloor + 1$, and the moment map is injective on them. At an interior equal-budget point all weights and budgets are positive, so no inequality reduces this local dimension. $\square$

In particular, for $R \geq 3$ the old certificate is generically not optimal after an unequal perturbation. A linear weighted estimate is therefore insufficient even arbitrarily close to equal budgets. The next theorem quantifies its loss.

I.76.5 Curvature and the precise loss of a frozen certificate

Let Π_Z be orthogonal projection onto Z . At a point satisfying Theorem I.76.2, let A_Z denote the positive definite operator $\Pi_Z L|_Z$. For a budget direction $v \in \mathbb{R}^R$, write $H_v = \frac{d}{dt} L(\varepsilon + tv)|_{t=0}$.

Theorem I.76.6 (Exact first and second variations). *On every open region of positive budgets satisfying the strict matrix certificate,*

$$DV[v] = -\frac{w^T H_v w}{1 - c} = -2 \sum_{j=1}^R f_j v_j, \quad (\text{I.76.13})$$

$$D^2V[v, v] = \frac{2}{1 - c} (\Pi_Z H_v w)^T A_Z^{-1} (\Pi_Z H_v w) \geq 0. \quad (\text{I.76.14})$$

At every strict-certificate point the Hessian has rank exactly $\lfloor (R-1)/2 \rfloor$. At an interior equal-budget point its kernel is the tangent space of the fixed-certificate polytope in Theorem I.76.5. At a zero budget, the same formulas describe the analytic extension of the rational branch, or its feasible one-sided restrictions.

Proof. The scalar $\kappa = -(1 - c)V$ is the minimum of $u^T L u$ under $\mathbf{e}^T u = 1$. The constrained quadratic is strictly convex on its affine domain, and its unique minimizer is w . Differentiation gives $\kappa' = w^T H_v w$. Differentiating the stationarity equation on Z gives $w' = -A_Z^{-1} \Pi_Z H_v w$. Consequently $\kappa'' = -2(\Pi_Z H_v w)^T A_Z^{-1} (\Pi_Z H_v w)$, proving the formulas.

Both L and H_v commute with reversal of coordinates, and w is reversal-invariant. Hence $\Pi_Z H_v w$ belongs to the symmetric subspace of Z , which has dimension $\lfloor (R-1)/2 \rfloor$. This image has full dimension. Indeed, if a symmetric $u \in Z$ is orthogonal to every $H_v w$, all nonconstant Fourier coefficients of

$$(c - \cos \theta) \Re(P_u(e^{i\theta}) \overline{P_w(e^{i\theta})})$$

vanish. The polynomial is constant and vanishes at $\theta = a$, so is identically zero. Symmetry of the coefficient vectors makes the product inside the real part real on the circle after cancellation of their common phase. Since P_w is not identically zero, $P_u = 0$, hence $u = 0$. This proves the rank assertion.

Finally $\Pi_Z H_v w^0 = 0$ is precisely the condition that w^0 remain stationary for the entire affine family $L + tH_v$. Along such a direction the same certificate and its positive representing weights persist locally. The dimensions in Theorem I.76.5 agree, proving the asserted identification of the kernel. $\square$

I.76.6 A rational example with a strict improvement

For $R = 3$, put

$$\ell_0 = c + \varepsilon_1, \quad \ell_1 = -c\varepsilon_1 - (1 - \varepsilon_2)/2, \quad \ell_2 = -c\varepsilon_2 + (\varepsilon_1 + \varepsilon_3)/2, \quad A = 6\ell_0 - 8\ell_1 + 2\ell_2.$$

The normalized symmetric minimizer has the form $(x, 1 - 2x, x)$, where $x = 2(\ell_0 - \ell_1)/A$. Whenever the strict certificate applies,

$$V_{3,a}(\varepsilon) = \frac{4(\ell_0 - \ell_1)^2/A - \ell_0}{1 - c}. \quad (\text{I.76.15})$$

This follows by minimizing $\ell_0 - 4(\ell_0 - \ell_1)x + Ax^2$.

Corollary I.76.7 (An explicit unequal third-frequency budget). *Let $a = \pi/3$ and $\varepsilon = (1/12, 1/12, 1/12 + t)$, with $|t| \leq 1/100$. Then*

$$\begin{aligned} V_{3,\pi/3}(\varepsilon) &= \frac{338}{3(91 + 12t)} - \frac{7}{6} \\ &= \frac{1}{14} - \frac{8t}{49} + \frac{96t^2}{49(91 + 12t)}. \end{aligned} \quad (\text{I.76.16})$$

The fixed equal-budget certificate gives only the first two terms. The displayed positive rational remainder is its exact loss.

Proof. Here $w = (x, 1 - 2x, x)$ with $x = 26/(91 + 12t)$, and (I.76.15) gives the formula. At $t = 0$ the restriction of L to Z is bounded below by $13/24$, while the perturbation of L has operator norm $|t|/2$. The matrix $M - sJ_4$ has smallest eigenvalue $13/28$ at zero. Its perturbation has norm at most $|t| + 4|s(t) - s(0)| < 2|t|$, because $|s'(t)| < 1/5$ on the specified interval. It stays positive. The value $s(t)$ is between 0 and 1. Finally $0 < x < 2/5$, and the three nonconstant coefficients before division by $1 - c$ are

$$\frac{3x^2 + (1 - 2x)^2}{2} - x(1 - 2x), \quad x(1 - \frac{5}{2}x), \quad \frac{1}{2}x^2.$$

They are all positive. The restriction of L to Z is positive and its constrained minimum is negative, which also verifies invertibility and $\beta < 0$. Theorem I.76.2 applies. $\square$

[write] **7 October 2026; sources 60–69.** *Checked at the write* (exact rational functions of t): $338/(3(91 + 12t)) - 7/6 = 1/14 - 8t/49 + 96t^2/(49(91 + 12t))$, and at $t = 0$ the value is source 37's $(\rho - e)/(1 + \rho)$ at $R = 3$, $a = \pi/3$, $e = 1/12$.

I.76.7 A quadratic finite-grid approximation

The second question in source 37 asks how much is lost when contact atoms must lie on a cyclic grid. The following theorem separates the endpoint displacement from the interior rounding error, and proves the proposed second-order upper approximation. It does not claim a complete exact solution of every finite grid problem.




Figure I.76.1: The exact gain over the old fixed certificate in the three-frequency example, as the third budget changes by t . The plotted quantity is $96t^2/[49(91+12t)]$; the proof gives this expression exactly. The vertical scale is in millionths.

Theorem I.76.8 (Aligned-endpoint grid bound). *Fix R, a in the strict window (I.76.7). On the angular grid $2\pi\mathbb{Z}/N$, let a_N be its smallest nonnegative allowed angle of size at least a , for the symmetric open exclusion I_a . Suppose $a_N < 2\pi/(R+1)$, and let ρ_R denote (I.76.8). Then*

$$0 \leq \inf_{\substack{\mu \text{ on } 2\pi\mathbb{Z}/N \\ \mu(I_a)=0}} \max_{1 \leq j \leq R} |\hat{\mu}(j)| - \rho_R(a_N) \leq \frac{\pi^2 R^2}{2N^2}. \quad (\text{I.76.17})$$

Uniformly for a in a compact subinterval of the strict window,

$$\inf_{\substack{\mu \text{ on } 2\pi\mathbb{Z}/N \\ \mu(I_a)=0}} \max_{1 \leq j \leq R} |\hat{\mu}(j)| = \rho_R(a) + \rho'_R(a)(a_N - a) + O_{R,a}(N^{-2}). \quad (\text{I.76.18})$$

In particular, when the endpoints align, the excess is $O_R(N^{-2})$.

Proof. The allowed grid is contained in K_{a_N} , so the continuous optimum $\rho_R(a_N)$ is a lower bound. Take its extremal measure. Its endpoint atoms are already on the grid. Replace each other atom, of mass u at angle θ , by the convex combination of its two adjacent grid angles with mean θ . Both neighbours are in the allowed complementary arc, using an unwrapped interval from a_N to $2\pi - a_N$.

For a twice differentiable vector-valued function g on an interval $[b, b+h]$, its linear interpolant at θ differs from $g(\theta)$ by at most $\frac{1}{2} \sup \|g''\|(\theta - b)(b + h - \theta)$. Apply this to $g(\theta) = e^{-ij\theta}$ and $h = 2\pi/N$. Each Fourier moment changes by at most $j^2 h^2/8$ after summing the atom masses. This proves (I.76.17); Taylor expansion of the explicit analytic function ρ_R gives (I.76.18). $\square$

For a half-open excluded interval, the two first allowed endpoints need not be symmetric. Formula (I.76.17) specifically uses the symmetric open convention. The same interpolation proof gives the quadratic rounding bound whenever the two endpoints of the chosen continuous comparison arc belong to the allowed grid; a shifted, possibly narrower comparison arc must be chosen first. No endpoint is silently moved into an excluded half-open boundary.

The practical implication for Lemma 5.2 is that unequal spectral information can be retained through a sharp finite-dimensional optimization, and discretization need not incur a first-order interior loss. Propagating the resulting constants through polynomial recurrence and a full Szemerédi iteration is a separate calculation.

Part II

The inverse step: energies, moments, restriction and dependent random choice

Contents. From many correlated derivatives to an approximate homomorphism. Sections II.1–II.4: source 02, delivered Sections 2–5 (a sharp weighted, mixed-function extension of Proposition 6.1 on every finite abelian group, its equality cases, deficits and rigidity). Sections II.5–II.6: source 03, delivered Sections 7–8 (a cubic dependent-random-choice certificate for Lemma 7.4, and an elementary replacement for Proposition 7.3). Sections II.7–II.10: source 04, delivered Sections 2–5 (its notation; exact second moments for Lemma 9.2; a polynomial restriction theorem for Lemma 9.3 and Corollary 9.4 over prime fields; arrangement moments for Lemmas 12.4, 14.7 and 14.8). The three sources treat different steps; no statement is printed twice. Source 04’s Section II.7 also fixes the cube notation used by its Part III sections.

Reading conventions of this Part (the full dictionary is Section F.3).

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
$\mathcal{E}$	02: $\mathcal{E}(w)$, additive energy of a weight on $G \times \widehat{G}$, normalized by N^{-3} ; 04: $\mathcal{E}_j(w) = M^{-1} \sum_{\chi} \widehat{w}(\chi) ^{2j}$ with the <i>unnormalized</i> transform, counting $2j$ -tuples; 03: $E(A)$ the additive energy of a set. <i>These are three different normalizations.</i>
$\mathcal{U}, T$	02: $\mathcal{U}(f) = \ f\ _{U^3}^8$ (normalized), $\mathcal{T}_w(f, g)$ a weighted derivative correlation, T its graph case; 04: $T_d(B)$ and $T_d^\phi(B)$ counts of additive and respected $2d$ -tuples.
κ, ρ, m	02: $\kappa = \mathcal{U}(f)/m(f)^4$, $m(f) = \mathbb{E} f ^2$; 03: $\kappa = 7/(3\sqrt{6})$, κ_* , $\rho = c/(2 - c)$, $m = V $ or $ A $; 04: $\kappa_d = 1/(2(400d)^d)$, $\rho = m/M$ with $m = A $.
R, L, Q	02: $R = \mathbb{E} K_u(s) ^2$; 04: R_j counts respected arrangements, L the kernel bandwidth, Q the kernel ratio sum (R14); 03: $L = D $. Source 04’s tags (M1)–(M7), (R1)–(R32), (A1)–(A14) label displays: <i>(R3) is a display, not a remainder R_3; (A1) is not Appendix A.</i>

Added in the second write (6 October 2026): sources 06–10. Contents. Sections II.11–II.19 are source 07’s chapter: a guide written for this report (Section II.11), source 07’s own notation table (its Appendix A, placed first), and its delivered Sections 2–8 (finite fibered domains and the directional error law; transport profiles; the coupled-path rigidity theorem; the sharp balanced threshold $1/(2k)$; the application to Gowers’s Section 10 with Freiman orders five and nine; the rational linear program for the starting law; collision-to-mode exactification, weighted laws and nonabelian groups). Sections II.20–II.21: source 10, delivered Sections 7–8 (a weighted energy extraction with $(c/4, 2^{10}c^{-5})$ in Proposition 7.3; the vertical-obstruction-fibre form of the Freiman extraction of Lemma 7.5, the order-eight retention $2^{-209}\gamma^{103}$ in Corollary 7.6, and a torsion budget for general finite codomains).

Overlaps (dated notes). Source 10’s Theorem II.20.1 and source 03’s Theorem II.6.3 are two elementary replacements of Proposition 7.3 under the same hypothesis; source 10’s is stronger in both conclusions (note after source 03’s theorem). Source 07 treats a step that no other source treats (Lemma 10.9); its relation to the Lean statement `lemma_10_9` is stated precisely in Section II.11 and in the note after Corollary II.17.6.

Reading conventions for the sections of sources 06–10 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
δ, δ_d, b_d	07: δ_d the directional defect of a direction d , b_d its boundary mass, δ a uniform bound on the δ_d . 07's δ is a failure probability, not a density.
$\mathcal{T}_s, \mathcal{C}, c(s)$	07: $\mathcal{T}_s(t)$ the transport profile, $\mathcal{C}(\mathbf{d})$ the cost of a word, $c(s)$, $A(s)$, $M(s)$ the coarse envelopes, ν an alternative starting law. 07's $\mathcal{C}$ is not the cube count $\mathcal{C}_d$ of 08 and 09 (Part III).
$\kappa, \tau, L, r_0, S, W$	07: κ the fibre-size ratio, L the one-step fibre-size variation, r_0 the internal lower bound, $\tau = L/r_0$; S the usable support, W the union of its fibres. 07's κ is none of the κ 's of 02, 03 and 04.
$\theta, \eta, \alpha, \rho, \sigma, M$	07: the parameters of Gowers's Lemmas 10.6–10.9 and of Section10.lean: $\theta = 10\eta^{1/5}$, $q(\theta) = 2\theta - \theta^2$, $\Delta = q(\theta) + \eta$; α, ρ, σ, M in (II.17.10) as there (M is the fibre-size cap, not a modulus). $\text{Fr}_k(D, H)$: Freiman homomorphisms of order k , the definition of the Lean <code>FreimanHom</code> .
$E(A), c, T(a, b)$	10: $E(A) = \sum_d r_A(d)^2$ and c as in 03's Theorem II.6.3; $w(a, b)$, $T(a, b)$ weighted codegrees, K a random subset, Z its bad pairs. 10's T is not 02's T .
$\Gamma, H_k, h_k, \gamma, F, V$	10: Γ a graph, H_k its vertical obstruction fibre and $h_k = H_k $, C a doubling constant, γ the graph-energy proportion, F the defect set and V the codomain of Theorem II.21.3. 10's h_k is not the $h_k = (k - 2)2^{k-1} + 1$ of Part IV, and this γ is not 10's γ of Part I.

Added in the third write (6 October 2026): sources 11–14. Contents. Sections II.22–II.32 are source 11's chapter: a guide written for this report and its delivered Sections 2–11 (graph energy and the exact endpoint; exact counting for one-valued errors; error-value comparisons for odd targets and targets with involutions; near-extremizers; correction without proximity; the inverse series; the exact largest non-affine energy for $|G| \geq 34$ and for all vector-space sizes; partial domains and Proposition 6.1). Sections II.33–II.42 are source 12's chapter: a guide and its delivered Sections 2–9 and Appendix A (translation loss along trees; sharp Freiman extension and its boundary; directional repair with the sharp majority lemma; finite Bohr-radius selection; fourth-moment smoothing; synchronization; the Section 10 local model with a polynomial radius; consequences; the fixed-domain repair of Lemma 10.6). Sections II.43–II.44: source 14, delivered Sections 4–5 (near-maximal energy and cube collisions force a unique nearby product coset with error $\tau(\varepsilon)$, sharp; equality with labels; equality and stability in arrangement moment interpolation).

Overlaps (dated notes). Collision-to-mode: source 07's proposition, source 11's Lemma II.28.1, source 12's sharp Lemma II.36.1. Section 10: source 07 treats Lemma 10.9, source 12 Lemmas 10.10–10.14; no statement is shared. Energy stability: source 11's near-extremizer bounds compose with source 14's Theorem II.43.2 (Remark W3, answering source 11's Question 3 in substance). Source 14's Section 5 refines source 04's arrangement interpolation (Theorem II.10.1). Source 11 is complementary to source 02's extension of Proposition 6.1. Remark W2 records that source 11's cutoff 34 cannot be dropped for general groups.

Reading conventions for the sections of sources 11–14 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
$\mathcal{E}, E, \varepsilon, d, t$	11: $\mathcal{E}(f)$ the ordered graph energy of $f : G \rightarrow W$, $E(S)$ the set energy of the error support, $\varepsilon(f) = 1 - \mathcal{E}(f)/N^3$, $d(f)$ the Hamming distance to affine maps, t an error density, Ψ_κ, D_κ . <i>11's $\mathcal{E}$ is not 02's normalized energy $\mathcal{E}(w)$.</i>
κ, A, C, J, M, T	11: $\kappa \in \{0, 1\}$ (involution in W or not); M, A, C, J, T counts of error values and triples (<i>not sets, not 07's or 02's T</i>); η a slack; ρ a test failure probability.
$b_B, \varepsilon, \kappa, \chi$	12: $b_B(t)$ the normalized translation loss, ε a missing density, κ a loss bound (<i>none of the other κ's</i>), $\chi(\eta) = (1 - \sqrt{1 - 2\eta})/2$ the majority function (<i>not a character</i>), Q_t collision probabilities, E_t edge sets.
$\mathcal{B}, \rho, \varrho, d, \zeta$	12: $\mathcal{B}(\Gamma, r)$ closed Bohr sets, ρ an outer radius, ϱ the component parameter of Lemma 10.5 (<i>not a radius</i>), $d = \Gamma $, ζ the final radius, $\eta = 2^{-43}$, $\theta = 10\eta^{1/5}$, $q = \sqrt{\theta}$ (<i>not an order</i>), L a fibre cap, σ an invariance error.
Q_ϕ, E, τ, S, R_j	14: $Q_\phi(A)$ congruent-cube collisions, $E(A)$ the energy (= 11's $E(S)$), $\tau(\varepsilon) = (1 - \sqrt{1 - 4\varepsilon})/2$ (= 12's $\chi(2\varepsilon)$), $S = \sum_r B_r ^3$ the profile (<i>not 14's Part III</i> $S = \ f\ _{U^2}^4$), R_j arrangement counts (as 04's), Ξ, S_{card} geometric defects, M a common mass.

Added in the fourth write (6 October 2026): sources 15–20. Contents. Sections II.45–II.56 are source 15's chapter: a guide and its delivered Sections 2–12 (axis cubes and collision measures; exact moment amplification; the cubic fibre bound; the product property coupled with the collision count; the comparison with the formal statements; Cartesian equality and stability; spectral stability; the joint normal form; a dichotomy; vertex weights). Sections II.57–II.64 are source 17's chapter: a guide and its delivered Sections 2–8 (notation; moments without density loss; one Fejér filter; the restriction theorem with all exceptional tuples; the limitation of the scalar certificate; fixed characteristic; boundary cases). Sections II.65–II.71: a guide, source 18's delivered Section 2 (degenerate arrangements at Lemmas 15.4–15.5) and source 19's chapter on Lemma 16.10, its delivered Sections 2–5 and 8 (the finite common-sampling problem; the exact minimax; sharp asymptotics; recovering polynomial pieces from common slices; worked comparisons).

Overlaps (dated notes). Source 15 against source 04's Section II.10 (its (A5)–(A11) again) and source 14's Sections II.43–II.44 (equality, the joint normal form and stability again, the last two with weaker constants); new in source 15 are the collision-parametrized product step and the coupled exponents. Source 17 against source 04's Sections II.8–II.9 (moments again; a complementary restriction theorem: better constant, worse threshold). Source 18's Section 15 chapter continues the chain from Lemma 14.8 that source 04, source 14 and source 15 hold; source 19's Section 16 chapter is new ground. **Cross-source answers.** Source 15's first question is answered in its leading constant by source 14 (item II.31); source 17's Question 7 partly by source 10 (item II.36).

Reading conventions for the sections of sources 15–20 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading in italics</i>)
arrangements	15: $X = \prod_i G_i$, $M = X $ (<i>not 14's common mass</i>), T the vertical group (14 writes R), $n = T $, V the value group, $q = V $; B_x , b_x , β_x fibres; β , ρ_3 (= 04's σ); R_d , $R_d(h)$; $C = R_1$ (<i>not 14's coset</i>), C_h , C_x ; $\Gamma = \sum_x b_x^3$ (14's S); a_d , χ normalized counts (χ <i>is not a character</i>); $Q(A)$ (= 14's); $e_k = 2k4^k$, $u_k = 8 \cdot 2^k$; $\mathcal{D}_d$ a defect; τ a band width (<i>not 14's $\tau(\varepsilon)$</i>); PP_γ the product property.
restriction	17: $m = 2r$ the number of variables (<i>17's T_{16} is 04's T_8</i>), T_m , G_m (= 04's T_d^ϕ), ε the relation vector (04's v), n a coefficient vector (04's ε), $L = \lceil (m+1)/(\alpha\eta) \rceil$, $S_{m,L}$ (04's Q), q a filter mean (<i>not 04's bandwidth</i>), $a = L^{-m}$, D_0 ; " L -degenerate" (<i>not 18's degenerate arrangements</i>); γ relative to $ B ^3$ (<i>the formal Gamma A d d i t i v e is absolute</i> , γN^3).
Section 15	18: d the arrangement order, k the horizontal dimension, $q = 2d$, $D = (2d+1)k + 2d - 1$, η a certificate, η_0 parity, $z_k(p)$, $Q_d(p)$, A_n , $J_{k,d}$, C_d , T_n , $M = 2^{k+4}$, E_k , λ (<i>these A_n, C_d, E_k are none of the earlier ones</i>).
Section 16	19: J the column set, $n = J $, H rows, $C_t(h)$ classes, q classes per row, d the interpolation order, r anchors, σ a lost mass, $W_{n,r,d}$, $\mathcal{M}_{n,q,d}(r)$ the minimax value, C_d and λ Poisson quantities, φ the golden ratio (<i>not a map</i>), M a list size, e_j elementary symmetric polynomials, $\mathfrak{c}$, $\mathfrak{q}$ control functions.

Added in the fifth write (6 October 2026): sources 21–27. Contents. Sections II.72–II.80 are source 24's chapter: a guide and its delivered Sections 2–9 (a moment–defect selector; a three-sample Lemma 7.4; the sharp scalar relaxation; walks from a profile certificate; an energy-to-small-difference-set theorem; Plünnecke–Ruzsa with proof; vertical-fibre compression and an extraction for $\mathbb{F}_p^d$ codomains; compiling the bounds). Sections II.81–II.103: a guide, source 22's delivered Sections 2–5 (the degree of a certificate; an exact degree-one count; the stratified finite bound; the two-cube case), source 25's delivered Sections 2–8 (minimal Walsh layers; codimension two; the cross-section hyperplanes; the global count and its geometry; ambient-size savings; a selection theorem), source 27's delivered Sections 2–11 (Boolean coefficient extraction; exact direction counts; first order in every characteristic; second-order structure; the second coefficient for $d = 2$; extensions; an exact eight-vertex count; budgets; safeguards), and source 23's delivered Section 8 (common base partitions and exact transport at Lemma 16.10).

Overlaps (dated notes). Source 24 against source 03's Sections II.5–II.6 (the cubic certificate, the scalar optimum and the two-witness barrier again) and source 10's Sections II.20–II.21 (the fibre extraction and the modern-input corollary again, with weaker or equal constants); new in source 24 are the extraction for $\mathbb{F}_p^d$ and the parameterized selector. Sources 22, 25 and 27 against source 18's Section II.66: the leading coefficient and the sharp single-certificate bound are source 18's, re-derived by 25 and 27 without citing it; source 22's $d = 1$ theorem contains source 27's Proposition II.102.1. Source 23's sampling lemma against source 19's Theorem II.67.4. **Cross-source answers.** Source 18's first question is answered for $d = 2$ by source 27 and its question on $d = 1$ in the leading coefficient by source 22 (item II.38); source 19's Question 6 is advanced by source 23 (item II.40).

Reading conventions for the sections of sources 21–27 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Section 7	24: μ an index measure, $c(i, j)$ codegrees, β a codegree threshold, η a defect penalty, M_t , B_t , Δ the moment–defect quantities, z_t , C_t the scalar root and constant ($C_3 = 03$'s κ_*^2), $\kappa = E(A)/m^3$ ($= 10$'s γ , 03 's c), τ a popularity threshold, D popular differences, p an edge density (<i>not a prime</i> in its Sections 5–6; a prime in Section 8), $Q_k = F_k $ (10 's $h_k = H_k $), L an avoidance length ($= 10$'s L).
Section 15	22: $q = 2d$, $D = (2d + 1)k + 2d - 1$, η a certificate, t its degree, $Q_t(p)$, A_n ($= 18$'s A_n), f_k , b_k , $J_{k,d}$ (<i>not 18's $J_{k,d}$</i>), U_t , C_d ($= 18$'s), $\mathcal{P}_{k,d}$, $\mathcal{A}_{k,d}$ thresholds, $L_k(p)$, R_k for $d = 1$. 25: $n = 2d2^k$, H_d ($= 18$'s C_d), $L_{k,d} = k + H_d$, $U_{k,d}$, $e_d(q)$ ($= 18$'s Q_d), $\mathcal{W}$, $\mathcal{H}_d$, $b_{k,d}(q)$, a_n , b_n , R_n , K_n , t . 27: $n = 2d$, $m = n2^k$ (<i>25's n is 27's m</i>), $\Delta_{k,d}(q)$, L_d ($= 18$'s C_d), $L_{d,p}$, $R_{k,d}$, $S_{k,d}$, T_r , B_ℓ , Q_ℓ , U_ℓ , U the cross-section space (25 's $\mathcal{W}$), s ($= 25$'s w , 18 's s), G , $g_{d,p}(q)$, $b_{2,d,p}$.
Section 16	23: proper k -boxes over $\mathbb{F}_p$ with one common step, m a width, $B_m = \lfloor m/2 \rfloor + 1$, $L_m = \lfloor (2B_m + 1)/3 \rfloor$ (<i>not 22's $L_k(p)$</i>), a_i , ε_i , Q_i slice-cover parameters, $A = \prod_i a_i$, $\eta = \sum_i \varepsilon_i$, x_i anchors, r their number, q fibre classes (19 's q), δ a sampling loss (19 's σ), $D_k = 2^{2^{k+8}}$ (<i>not a dimension</i>).

Added in the sixth write (6 October 2026): sources 28–33. Contents. Sections II.104–II.113: a guide, source 31's delivered Sections 2–5 and 9 (normalized derivatives and their phase space; an extremal subgroup without exponent loss; the sharp local quartic envelope; sharp stability of a selected derivative graph; the stability statement for the weights) and source 29's Sections 5–8 (joint endpoint rigidity; soft selectors; mixed selectors; sharpness of the error order). Sections II.114–II.118: a guide and source 31's Sections 6–8 and Appendix A (weighted energy extraction; a weighted Freiman restriction; keeping the derivative correlation; a square-root weighted mass theorem). Sections II.119–II.120: a guide and source 29's Section 9 (restriction without discarding affine degeneracies). Sections II.121–II.135: a guide, source 28's Sections 2–7 (conventions; weighted local errors and an anchor; logarithmic selection; exact local differences; Fourier smoothing and the main theorem; consequences) and source 30's Sections 2–9 (translation calculus; the escape-tube inequality; packing; dense-domain extension; Corollary 10.11; a compatibility step for Lemma 10.12; consequences; exponential-rank obstructions). Sections II.136–II.143: a guide and source 33's Sections 2–8 (sampling and normalization; a noisy homomorphism lemma; synchronizing derivatives; sharp local defect inequalities; extremal families; erasures; exact cyclic decoding).

Overlaps (dated notes). Source 29's endpoint constants are superseded by source 31's; source 31's restriction is source 10's with weights, with weaker order-eight constants and source 10's Reiher–Schoen formula; source 30's tree budget and compatibility step are source 12's, and its extension budget a weaker form of source 12's; source 28 credits sources 07 and 12 for paths, extension and smoothing. **Cross-source answers.** Source 02's Subsection 11.1 is answered in a prime-cyclic, bounded-fibre form by source 31 (item II.1) and its Subsection 11.2 advanced by sources 31 and 29 (item II.2); source 04's Question 4 and source 17's Question 2 are answered for fixed fields by source 29 (item II.6); source 12's Questions 4–6 are advanced by source 30 (item II.28); the label-repair questions of sources 11, 14 and 15 are advanced on full products by source 33 (item II.24).

Reading conventions for the sections of sources 28–33 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Proposition 6.1	31: $V = G \times \widehat{G}$, W_z Weyl operators, $q = q_f$ the normalized squared derivative coefficients ($0 \leq q \leq 1$, $\sum_z q(z) = N$; <i>not 29's field size</i>), $\kappa = \ f\ _{U^3}^8 / \ f\ _2^8$, $\delta = 1 - \kappa$, F the best squared correlation with a coset phase, $t = 1 - F$, $\mathcal{R}(\delta)$, $\mathcal{G}(\varepsilon)$, $\mathcal{Q}_{\text{cos}}(G)$ coset phases, H an isotropic subgroup, $\mathcal{E}(w)$ and $\mathbf{E}$ weighted energies (<i>two different letters</i>), Θ_w a weighted score, T_ϕ the normalized graph score. 29: $A_f(h, \chi)$, $T_\phi(f)$ (unnormalized by m^2 ; 31's T_ϕ is $T_\phi(f)/m^2$), $m = \mathbb{E} f ^2$ (<i>not 29's moment index $m = 2^{d-2}$ in Part III</i>), ψ a derivative frequency map, ε the endpoint defect.
Section 7	31: w weights, M the mass, c the relative energy ($= 10$'s γ), K a fibre cap (<i>not a frequency set</i>), V_k , $q = V_k $ ($= 10$'s H_k , h_k), C a doubling constant, μ, ν, b spectral profile statistics, τ a threshold.
Section 9	29: $\mathbb{F}_q$ a fixed field, $\mathcal{R}$ the relation space, $d = \dim \mathcal{R}$, $s = \ell - d$ free coordinates (<i>not 33's s</i>), $r(x)$, $e(x)$ ranks, $\lambda = q^{-t}$, $S = S_{q,s}$ a rank sum.
Section 10	28: (X, r) a domain over G , X_s fibres, $g(s) = X_s $, Γ the original-domain spectrum, Bohr(Γ, u), $p(\alpha)$ the inclusion probability (<i>not a prime</i>), $\lambda(\alpha)$, R_α , ζ_α , η a defect, T an anchored energy, L a fibre bound. 30: $B(\chi, \rho)$ rectangular Bohr sets with widths $\rho_i \leq 1/4$, $b_B(d) = B - B \cap (B + d) $ (a count; 12's b_B is a fraction), m a step multiplier, $c_r = 3^r - 1$, D_B a measured growth ratio, κ ($= 12$'s κ), η ($= 12$'s ε), $\mathcal{L}_r$, $S_r(\varepsilon)$ extremal coefficients.
Section 12	33: $s = 2m$ vertical differences ($s = 16$ is Gowers's <i>eight-arrangement</i>), ε_s the failure probability, $\mathcal{V}(G, H; K)$ the models $a(x) + B(x, y) + \ell(y)$, δ the distance (<i>not a density</i>), R_s the repair modulus, F_s , Q_s local polynomials, t_h a derivative profile, θ (Lemma 12.3's parameter).

Added in the seventh write (6 October 2026): sources 34–39. Contents. Sections II.144–II.148: a guide and source 35's Sections 2–5 (the arrangement model and the quantitative theorem; degree balances; the random restriction and its collision correction; the exponent and a downstream substitution). Sections II.149–II.150: a guide and source 38's Section 6 (the sharp odd-order local stability coefficients). Sections II.151–II.161: a guide and source 39's Sections 2–11 (a finite model of shared-anchor certification; a cutoff-free sampling theorem; the exact finite minimax; a closed profile formula and the affine cap; the Poisson profile; uniform asymptotics; deterministic selection; polynomial reconstruction and errors; labelled slice covers; label-preserving refinement). Sections II.162–II.163: a guide and source 34's Appendix A (the sharp nonuniform directional rigidity threshold).

Overlaps (dated notes). Source 39 re-derives source 19's sampling theory (sampler, minimax, permutation orbits, the constant $\varphi^3 e^{-\varphi}$) without citing it; source 35's one-stage selector and retained powers are source 40's (placed, not yet written); source 38's noncommuting-subgroup gap is source 31's. **Cross-source answers.** Source 07's Question 1 is answered by source 34 (item II.9); source 31's odd-order coefficients by source 38 (item II.46); source 19's Question 1 for source 39's rule, not for source 19's own (item II.39).

Reading conventions for the sections of sources 34–39 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Section 15	35: $m = 2d$ cubes of dimension k , $s = m2^k$ vertices, L the Fejér bandwidth (<i>not a cell length</i>), μ a product law with largest atom $1/L$, D_t the dimension of the degree- t balance space, $E_{k,m}$ the field-size exponent, $M = (m+1)k + m - 1$, $T_0 = \beta^{m-1}q^M$, $W = \sigma^\perp$, q the field size.
Proposition 6.1	38: v a unit vector, $\varepsilon = 1 - \ v\ _{U^3}^8$, $\mathcal{F}(v)$ the best coset-phase fidelity, $\mathfrak{S}_{\text{odd}}$, $\mathfrak{S}_{6'}$ moduli, W_h Weyl operators, B the commutator bicharacter, $\kappa(v) = \ v\ _{U^3}^8$ (= 31's κ), $\langle \cdot, \cdot \rangle$ as in source 31.
Lemma 16.10	39: J the m columns, q classes per row (= 19's q), t the anchor threshold (= 19's d), r anchors, $\ell_{m,r,t}(s)$ its one-class loss (<i>not 19's $W_{n,r,d}$, which discards sampled points of failed classes</i>), $\mathfrak{M}_{m,q,r,t}$ (= 19's $\mathcal{M}_{n,q,d}(r)$ only in the limit), F_t , λ_t , κ_t , φ , L list sizes, W a total row weight.
Section 10	34 (Appendix A): κ the fibre ratio (= 07's κ), δ_d failure probabilities, k the Freiman order, R_g fibre sizes, Ω their sum.

Added in the eighth write (7 October 2026): sources 40–46. Contents. Sections II.164–II.172: a guide and source 40's Sections 2–9 (arrangements and the counting normalization; a single monomial isolates the universal relation; bounded-coefficient exceptional arrangements; Fejér kernels and exact survival; the one-stage selection theorem; comparison with Lemma 15.5; vector targets and torsion; a sharp exponent obstruction for positive Fourier kernels). Sections II.173–II.182: a guide and source 44's Sections 2–10 (arrangements and the regular region; a width theorem; exact core-avoidance laws; the full degeneracy count; the second Whitney coefficient and a uniform cutoff; the nine-point grid; eight lattice energies and the $d = 8$ coefficient; the pair-overlap deficit; growing-dimension laws). Sections II.183–II.185: a guide and source 45's Sections 2–3 (the first three additive-energy levels; a larger coset-stability range and the geometry of equality).

Overlaps (dated notes). Source 35's degree-weighted selection improves source 40's field-size exponent; source 44's core count and $d = 2$ coefficient are source 27's; source 45's first energy gap is the inverse Pollard theorem's and its envelope source 14's on a smaller range. **Cross-source answers.** Source 27's two questions are answered by source 44 (item II.38); source 14's Question 4 is advanced by source 45 (item II.31); source 04's and source 17's retention questions gain source 40's barrier (item II.7).

Reading conventions for the sections of sources 40–46 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Section 15	40: $s = 2d2^k$, $M = (2d+1)k + 2d - 1$, $E_k = s2^{s-1}$, L the Fejér bandwidth, $\Gamma_s(\alpha, \eta)$ the retention, $a_{s,L}$, $b_{s,L}$, $R_s(L)$ good and bad survival and their ratio, $T(B)$, $T_0 = \beta^{2d-1}q^M$, $G_\phi(B)$, ρ an arrangement, η the accuracy, η_0 the universal relation.
energy, cosets	44: q the field size, p its characteristic, $D = (2d+1)k + 2d - 1$ (= 40's M), $m = 2d$, H_d (= 18's C_d), $b_{2,d}$ (= 27's $b_{2,d,p}$), $\{-1, 0, 1\} = \{-1, 0, 1\}$ (<i>not the circle</i>), $Z_{k,d}(q)$ the degenerate count, $S_{k,d}$ the core count.
	45: $E(A)$ the additive energy, $M_m = (2m^3 + m)/3$, $D(A) = (M_m - E(A))/4$ (<i>not 44's D</i>), H_m , J_m sets, $\varepsilon = 1 - E(A)/m^3$, $\tau(s)$ (= 14's τ), $\delta = A \triangle C /m$, $\eta = \varepsilon - \delta(1 - \delta)$, $r(d) = A \cap (A + d) $, H the canonical subgroup (<i>not 44's H_d</i>), K , D holes.

Added in the ninth write (7 October 2026): sources 47–53. Contents. Sections II.186–II.194: a guide and source 48's Sections 2–9 (sampling counts and the repeated-entry obstruction; an exact group-algebra inventory; local replacement and Stirling coefficients; deterministic signed-energy rounding; exact cardinality; general linear systems; robust certificates and rational filters; a deterministic polynomial-loss restriction theorem). Sections II.195–II.219: a guide, source 53's Sections 2–8 (the context and the quantifiers; scalar parameters; seven budgets; the exact threshold collapse; the short witness; short parents and square geometry; final assembly), source 51's Sections 2–10 (the retained interface; the obstruction; uniform constants;

a strip context and its initial data; the six budgets; the output chain; failure of the grid; a conditional covering repair; a direct repair) and source 52's Sections 2–9 (the interface; the dyadic strip; the step-two counterexample; genuine initial stages; unit input steps; all six budgets; prime-power families; a two-point model). Sections II.220–II.225: a guide and source 49's Sections 2–6 (two-cube degeneracy as a translation problem; rank two and the second coefficient; exactness in dimension three; eventual polynomiality; finite-field extensions and the subfield alphabet). Sections II.226–II.235: a guide, source 50's Sections 2–9 (definitions; one-dimensional cover laws; product boxes and global rigidity; balanced words; profile realization; the pinned Section 16 interface; the counterexample; a repair interface) and source 49's Section 8 (the acyclic cover theorem).

Overlaps (dated notes). Source 48's restriction is source 17's, made constructive; source 53's geometric lemmas are source 51's; source 49's first-order count and translation lemma are source 22's, its closing-comparison audit source 19's. **Cross-source answers.** Source 17's Question 9 is answered in ambient size by source 48 (item II.23); source 22's exactness question for $k \leq 3$ by source 49 (item II.42); source 19's and source 23's search for controls closing Lemma 16.10 is answered negatively for the catalogue form by source 50 (item II.41).

Reading conventions for the sections of sources 47–53 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Section 9	48: s the order, $m = 2s$, n the number of labels, D the total group size, $N_{a,b}$ partial energy counts, $E_s(S;g)$ (<i>not 45's $E(A)$</i>), $C_s = \binom{s+2}{2}^2$, $H_p(t)$, $R_{p,q}(t)$ response series, $\mathcal{T}_s$, $\mathcal{R}_s$ total and respected tuple counts, L the Fejér bandwidth (= 17's), $d = \alpha + \eta - \alpha\eta$ a slack (<i>not a dimension</i>), Q a rational grid (<i>not 50's or 51–53's Q</i>).
Section 13	51–53: a or α the actual density, a_0 its lower bound, K , Q , c , β , f , g , γ the scalar functions (= the formal section13Q , section13SquareExponent), θ the Section 10 cutoff, t , b Stage 13.4 bounds, u , v , w recurrence exponents, z the rounded Bohr radius, $W = 2^{135}a^{-704}$, $D(a)$ a coefficient (<i>not a set</i>), m the Stage 13.6 integer, P , Q_{st} (or Q), R , S , U the stage progressions, $\mathcal{H}$ the critical heights; 51: $\ell(P)$, d_P , $K = 2^{762}$, $T = 7,995,833$; 52: ℓ , $\ r\ _N$, T the strip, K_0 , $q = 2(2K_0 + 1)$.
Lemma 15.4	49: k the cube dimension, $\delta_{k,1}(p)$, $L_k(p)$ (= 22's), $\mathcal{P}_k(p)$ ternary multiaffine polynomials, $T_k(p)$ the translation locus, $\mathcal{B}_k(p)$ the signed-coordinate arrangement, $r(f)$ derivative rank, M_k the rank- ≥ 3 count, χ_k a characteristic polynomial ($d = 1$ <i>here; not 27's or 44's cross-section convention</i>).
Lemma 16.10	50: R the alphabet size, $b_{(j)}$ ordered symbol masses, B_s , $M_s(\pi)$ branch profiles, $\text{Cov}_{q,d}$ the list-cover optimum, $\text{ML}_{\gamma,r}$, E_D , $c(\cdot)$, $q(\cdot)$, $s(\cdot)$ the catalogue controls, Q_γ , Q_δ graph counts, r_0 ; 49: N_j , Q_j , $c_{j,D}$, $T_{j,D}$ the four controls, II–I7 the interfaces.

Added in the tenth write (7 October 2026): sources 54–59. Contents. Sections II.236–II.247: a guide, source 55's Sections 2–11 (main results; demodulation; four cube vertices; a cubic remainder; a two-torsion profile; the explicit local theorem; uniform entry and extremizing families; rigidity; the exact U^2 calibration; consequences) and source 57's Section 5 (near-maximal uniformity and global frequency maps). Sections II.248–II.260: a guide, source 59's Sections 7–10 (signed graph expansions; averaging the signed kernel, radii 4/9 and 1/2; classical correction and global recovery; examples, comparisons and claim boundaries) and source 56's Sections 2–9 (normalization; autocorrelation and triangle variance; weighted modal differences; self-correction and coarse repair; the one-error profile; the main proof; exact extension and testing; quadratic level sets). Sections II.261–II.266: a guide and source 59's Sections 2–6 (exact character distributions; adaptive separation; weighted Freiman restrictions over arbitrary targets; small doubling and Lemma 7.5; one-character restrictions and lonely runners). Sections II.267–II.275: a guide and source 54's Sections 2–9 (the mixed weighted Fourier factorization; quadratic untwisting; every finite abelian group; offset ceilings; stability and the fourth-root obstruction; several intercepts; localization; the deficit and a selector). Sections II.276–II.286: a guide and source 58's Sections 2–11 (counts; admissible filters; the random selector; nonprincipal relations; repeated bad tuples; the expectation theorem; powered cosines and the exponent 54;

generalizations; the optimal leading exponent; moment lifting and the interfaces). Sections II.287–II.290: a guide and source 57’s Sections 2–4 (the imported results; Fejér purification; propagation and the printed exponent). Sections II.291–II.292: a guide and source 57’s Section 6 (bounded Schur defects and the third energy level).

Overlaps (dated notes). Source 55’s $d = 3$ coefficients are those of sources 38 and 31; source 59’s affine cubic is source 11’s; source 58’s and source 48’s restrictions both count repeated entries; source 57 reproves source 40’s Fejér selection in the needed specialization and source 45’s energy recurrence. **Cross-source answers.** Source 45’s first two questions are answered by source 57 (item II.53); source 17’s composite-modulus question by source 58 (item II.8); source 24’s mixed-torsion question by source 59 (same item); source 11’s cutoff question is advanced by source 59 (item II.21); source 57’s threshold question is answered by source 53 (item II.55).

Reading conventions for the sections of sources 54–59 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Prop. 6.1 endpoint	55: $q = 2^d$, $\mathcal{U}_d(f) = \ f\ _{U^d}^q$, $\varepsilon = 1 - \mathcal{U}_d(f)$, $D_d(f)$ the squared distance to $\mathcal{P}_{d-1}(G)$ (<i>not 56’s $\mathcal{D}_w$</i>), $\mathfrak{S}_{d,G}(e)$ the modulus, m, r, a, D, z, v the demodulated quantities, C_d ($= 04$ ’s P_d), A_d, J_d, B_d cube-vertex counts, $Q_d(r)$ a quartic circuit sum (<i>not 57’s $Q_d(f)$</i>); 57: $Q_d(f) = \ f\ _{U^d}^{2^d}$, Δ_h forward derivatives, C_s, K_k correction constants.
Affine rigidity	59: $\varepsilon(f) = 1 - \mathcal{E}(f)/N^3$, t the error density, $\Psi(t) = 4t - 10t^2 + 6t^3$ ($= 11$ ’s Ψ_0), $D(u)$ its inverse on $[0, 1/7]$, a, d colour and support deficits, $W[r]$ the r -torsion; 56: w a weight of mean α , $u = \ w - \alpha\ _{U^2}$, $\lambda = u^2/\alpha^{5/2}$, $\kappa = u^4/\alpha^4$, $s = u^2/\alpha^2$ (<i>not 59’s d or 57’s s</i>), $\mathcal{Q}(w)$, $\mathcal{D}_w(\phi)$, $\varepsilon = \mathcal{D}_w/\mathcal{Q}$, t_w the relative distance.
Section 7	59: q the number of sign orbits of a forbidden set F , $\ \cdot\ _{\mathbb{T}}$, F_k^0 the order- k vertical defects, $\beta_n(\theta)$ the small-ball law (<i>not 57’s density β</i>), $m_b(q)$ the ranks, C the doubling constant.
Section 8	54: $\sigma = \delta(1 - \delta)$, $B = \max(\delta, 1 - \delta)$, $D_h(r)$ Gowers’s coefficient, $A_\eta(g)$, $\mathcal{E}_\eta(g)$ the offset energy, $\mathcal{D}_\eta$ the spectral overlap, $\mathcal{W}_{\xi,w}$, $\mathcal{L}_{\xi,w}$ windows and scores, c_m the offset ceilings, K_m stability constants, p_ξ normalized powers (<i>not a prime</i>).
Section 9	58: $\mathsf{T}_m(B)$, $\mathsf{G}_m(B, \phi)$, U_m tuple counts (<i>not the report’s $\Gamma_{\delta,s} = \Gamma_{\delta,s}$</i>), c_n , P , Q , R_s filter quantities, D the filter degree, $\tau_G(L)$ the domain torsion, $\mathcal{A}_{m,a}(K)$ the principal cost, a_K , z_K its exponents.
Section 13, energies	57: $a = \alpha/2$, $b = 14,409,429$, $d = 2,996,201,429$, $\delta = a^d$, $\gamma(\delta)$, $\mathcal{Q}(\delta)$ ($=$ the formal section13SquareExponent , section13Q), $T(A)$, $R(A, \phi)$ arrangement counts, N_{sq} the inherited threshold; $h(C)$ the Schur deficit, $E(A)$, M_m , $D(A)$ ($= 45$ ’s), J_m , K_m the third-level sets (<i>not 57’s constants K_k</i>).

Added in the eleventh write (7 October 2026): sources 60–69. Contents. Sections II.293–II.295: a guide and source 67’s Sections 3–4 (cubic rigidity throughout the positive range in characteristic three; sharp thresholds for global affine recovery). Sections II.296–II.309: a guide, source 68’s Sections 2–4 (a sharp pointwise energy inequality; coset rigidity and two thresholds; puncture repair) and source 66’s Sections 2–11 (ordered defects and endpoint reflection; arithmetic prefixes; the partition normal form; boundary holes; the hole energy law; the endpoint envelope; the first and second levels; the third level; the fourth level; inverse consequences). Sections II.310–II.330: a guide, source 68’s Sections 5–7 (decaying exceptional terms; the degree-twelve filter and exponent fifty-three; uniform principal costs), source 62’s Sections 2–9 (the main theorem; the unit-dilation detector; primitive responses; weighted short relations; the restriction proof; the order-eight strengthening; torsion obstructions; extensions), source 61’s Sections 2–9 (the relation interface; separating defects; the vertical budget; exact partitions; energy to growth; exact restriction of high energy; consequences for Gowers’s lemmas; torsion obstructions) and source 60’s Section 9 (a translated Fejér restriction theorem). Sections II.331–II.337: a guide and source 60’s Sections 3–8 (respected arrangements; a spectral inequality with independent labels; the complete polynomial and its equality cases; repair to every order; odd targets and the common-prime interval; two-direction testing). Sections II.338–II.339: a guide and source 67’s Section 5 (the exact number of interpolation candidates).

Overlaps (dated notes). Source 67’s local range contains source 59’s $t < 4/9$ for its target class, and its order-thirteen cutoff uses source 59’s example; source 68’s restriction continues source 58’s and source 62’s precedes both in date; source 61 finds source 59’s exponent independently; source 66’s normal form and third level are source 57’s in another form, the same day; source 60’s arrangement polynomial extends source 33’s cubic. **Cross-source answers.** Source 45’s Questions 6.6–6.8 and source 14’s Question 4 are answered by source 68 (item II.31); source 57’s Research questions 2–3 are answered and its Research question 1 advanced by source 66 (item II.53); source 58’s Questions 13.1–13.2 are answered or advanced by source 68 (item II.7); source 33’s Questions 2, 3 and 8 are answered and its Question 1 advanced by source 60 (item II.49); the intake’s marked threshold composition is proved with a smaller threshold by source 60 (item II.6); source 11’s cutoff question is settled for elementary abelian 3-group targets by source 67 (item II.21).

Reading conventions for the sections of sources 60–69 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Characteristic three	67: W an elementary abelian 3-group, $\mathcal{E}(f)$, $\varepsilon = 1 - \mathcal{E}(f)/N^3$, t , $\Psi(t) = 4t - 10t^2 + 6t^3$ (= 59’s and 11’s), D its inverse, a, b, c colour and rank remainders, A, C, T quadruple counts; its thresholds $1/4$ and $4/9$ bound the <i>defect</i> ε (<i>not 59’s radius</i> $t < 4/9$).
Energies	68: $m = A $, $r_A(d) = A \cap (A + d) $, $q(d) = r_A(d)/m$, $E(A)$, $\varepsilon = 1 - E(A)/m^3$, $\tau(u) = (1 - \sqrt{1 - 4u})/2$ (= 14’s and 45’s), H, C the canonical subgroup and coset, δ , η , $s = \eta/\delta^3$, θ_d ; 66: $\sigma(C)$ the Schur defect (= 57’s $h(C)$), $e(B, a)$ the endpoint defect (<i>not the exponential</i>), $E(A)$, $D(A) = (M_m - E(A))/4$ (= 45’s E, D), $M_m = (2m^3 + m)/3$, $[u, v]_{\mathbb{Z}}$, $\text{wt}(H)$, $\kappa(H)$, $A_H(N)$, $p(r)$ partitions, $J_m^{\pm}$, C_m , $T_m^{\pm}$ templates (<i>C_m is 57’s K_m; 66’s $J_m^{\pm}$ are 57’s J_m in its two orientations</i>).
Restriction	68: $b = B $, T_{16} , G_{16} tuple counts (= 58’s T_{16} , G_{16}), $P, Q, R_s, \mathcal{R} = P/Q, \mathcal{A}$ and $\mathcal{I}_{m,a}$ principal costs (= 58’s $\mathcal{A}_{m,a}$), a a discrepancy order; 62: N, M the moduli, $\beta = B /N$, $m = 2k$, T_m, S_m counts, L the filter length, $H_{\text{el}}, H_{\text{RS}}$ envelopes, $\kappa_{N,m,L}$ the weighted degeneracy constant (<i>not 66’s $\kappa(H)$</i>), $\mathcal{R}_m$ the relation set; 61: Γ the graph, K its difference ratio, $\mathcal{D}_m(\Gamma)$ the vertical defects (<i>not 56’s $\mathcal{D}_w$</i>), $\mathcal{E}_s$ relation energies, $a_m = (2m + 1)\log_2(5m)$, b_m the colouring exponent (<i>not 68’s b</i>); 60: $L = \lceil (m + 1)/(\alpha\eta) \rceil$, $\tau_G(L)$ the domain torsion (= 58’s).
Arrangements	60: s the number of vertical differences, ε_s the failure probability, $\mathcal{V}, \mathcal{H}, \mathcal{A}$ the vertical, horizontal and biaffine models, δ the distance, $B_s(\delta)$ the polynomial, ρ_s its inverse branch, $E_s(t)$ the auxiliary profile, c_s^{prime} , $J(\eta)$, P_s, σ_s (33’s δ is this δ).
Lemma 16.10	67: r anchors, L_i list sizes, t the degree bound of fibre classes, E_t the elementary symmetric polynomial, T, D counts (<i>not 19’s or 39’s t, q</i>).

II.1 Normalization and finite Fourier identities

Source 02, delivered Section 2.

Let G be a finite abelian group of order $N \geq 1$, written additively, and let $\widehat{G}$ be its character group, written multiplicatively. Thus $(h, \chi) + (k, \psi)$ in $G \times \widehat{G}$ means $(h + k, \chi\psi)$. All averages over G are normalized; sums over $\widehat{G}$ are not:

$$\mathbb{E}_x a(x) = \frac{1}{N} \sum_{x \in G} a(x), \quad \widehat{a}(\chi) = \mathbb{E}_x a(x) \overline{\chi(x)}.$$

Fourier orthogonality gives

$$\mathbb{E}_x |a(x)|^2 = \sum_{\chi \in \widehat{G}} |\widehat{a}(\chi)|^2. \quad (\text{II.1.1})$$

For normalized convolution $(a * b)(s) = \mathbb{E}_h a(s-h)b(h)$, we have $\widehat{a * b} = \widehat{a} \widehat{b}$. These facts follow by expanding the sums and using $\mathbb{E}_x \chi(x) = 0$ for a nontrivial character and 1 for the trivial character.

Define

$$\Delta_h f(x) = f(x) \overline{f(x-h)}, \quad m(f) = \mathbb{E}_x |f(x)|^2,$$

and define the eighth power of the normalized U^3 norm by

$$\mathcal{U}(f) = \mathbb{E}_h \sum_{\chi \in \widehat{G}} |\widehat{\Delta_h f}(\chi)|^4. \quad (\text{II.1.2})$$

It also equals the usual cube average

$$\mathbb{E}_{x,a,b,c} \prod_{\omega \in \{0,1\}^3} \mathcal{C}^{|\omega|} f(x + \omega_1 a + \omega_2 b + \omega_3 c), \quad (\text{II.1.3})$$

where $\mathcal{C}$ denotes complex conjugation. One proof first expands the fourth Fourier moment in (II.1.2); character orthogonality leaves the additive quadruple constraint, which is parameterized by three variables. The remaining derivative variable gives (II.1.3), after changes of signs in the increment variables. In particular $\mathcal{U}(f)$ is nonnegative.

For $f, g : G \rightarrow \mathbb{C}$, write

$$A_{f,g}(h, \chi) = \mathbb{E}_x f(x) \overline{g(x-h) \chi(x)}.$$

This is just a Fourier transform of a mixed derivative, not an additional assumption about f or g .

Definition II.1.1 (Weighted correlation and energy). For any $w : G \times \widehat{G} \rightarrow [0, \infty)$, define

$$\mathcal{T}_w(f, g) = \mathbb{E}_h \sum_{\chi} w(h, \chi) |A_{f,g}(h, \chi)|^2, \quad (\text{II.1.4})$$

$$\mathcal{E}(w) = \frac{1}{N^3} \sum_{\substack{h_1+h_2=h_3+h_4 \\ \chi_1\chi_2=\chi_3\chi_4}} \prod_{j=1}^4 w(h_j, \chi_j). \quad (\text{II.1.5})$$

The factor N^{-3} in (II.1.5) is deliberate: it normalizes the energy of a full additive graph to 1. It is not the uniform-probability normalization on the group of order N^2 . If $w = \mathbf{1}_\Gamma$, then $N^3 \mathcal{E}(w)$ is the ordinary ordered additive-quadruple count of Γ .

II.2 A sharp weighted extension of Proposition 6.1

Source 02, delivered Section 3.

Theorem II.2.1 (Mixed energy-sensitive inverse step). *For every finite abelian group G , every $f, g : G \rightarrow \mathbb{C}$, and every nonnegative weight $w : G \times \widehat{G} \rightarrow [0, \infty)$,*

$$\boxed{\mathcal{T}_w(f, g)^4 \leq \min\{m(f)^4 \mathcal{U}(g), m(g)^4 \mathcal{U}(f)\} \mathcal{E}(w).} \quad (\text{II.2.1})$$

There is no pointwise boundedness assumption on f, g , and no graph or bounded-fiber assumption on w . The coefficient 1 cannot be decreased.

[write] **6 October 2026; sources 11–14.** Source 11's chapter (Section II.22) starts from the output of Proposition 6.1 in the range $\alpha^4 > 8/9$ and recovers an affine map (its Corollary II.32.2); it is complementary to this theorem, which strengthens the energy lower bound itself.

II.2.1 Three exact identities

For $u \in G$ define

$$F_u(s) = f(s)\overline{f(s-u)}, \quad D_u(s) = g(s)\overline{g(s-u)}, \quad V_u(h) = \sum_{\chi} w(h, \chi)\chi(u),$$

and put $K_u = D_u * V_u$. Use the normalized inner product

$$\langle F, K \rangle = \mathbb{E}_{u,s} F_u(s) \overline{K_u(s)}.$$

Then

$$\mathcal{T}_w(f, g) = \langle F, K \rangle, \tag{II.2.2}$$

$$R := \mathbb{E}_{u,s} |K_u(s)|^2 = \mathbb{E}_u \sum_{\psi} |\widehat{D}_u(\psi)|^2 |\widehat{V}_u(\psi)|^2, \tag{II.2.3}$$

$$\mathbb{E}_u \sum_{\psi} |\widehat{V}_u(\psi)|^4 = \mathcal{E}(w). \tag{II.2.4}$$

For (II.2.2), expand the square in (II.1.4), set the second spatial variable equal to $s - u$, and obtain

$$\mathbb{E}_{u,s} f(s) \overline{f(s-u)} \mathbb{E}_h \overline{g(s-h)} g(s-u-h) \sum_{\chi} w(h, \chi) \overline{\chi(u)}.$$

This is exactly $\langle F, K \rangle$. Nonnegativity of w also makes the original expression visibly real and nonnegative. Identity (II.2.3) is Parseval and the convolution identity.

For (II.2.4), expand

$$\widehat{V}_u(\psi) = \frac{1}{N} \sum_{h, \chi} w(h, \chi) \chi(u) \overline{\psi(h)}.$$

The fourth power contributes N^{-4} . Summing over ψ contributes N and imposes $h_1 + h_2 = h_3 + h_4$. Averaging over u imposes $\chi_1 \chi_2 = \chi_3 \chi_4$ without another factor. The remaining normalization is N^{-3} , proving the identity.

II.2.2 Proof of the theorem

The first exact norm is

$$\|F\|_2^2 = \mathbb{E}_{u,s} |f(s)|^2 |f(s-u)|^2 = m(f)^2.$$

Cauchy–Schwarz applied to (II.2.2) therefore gives

$$\mathcal{T}_w(f, g)^2 \leq m(f)^2 R. \tag{II.2.5}$$

Applying Cauchy–Schwarz to (II.2.3), this time on the measure space with average in u and counting measure in ψ , gives

$$R^2 \leq \left(\mathbb{E}_u \sum_{\psi} |\widehat{D}_u(\psi)|^4 \right) \left(\mathbb{E}_u \sum_{\psi} |\widehat{V}_u(\psi)|^4 \right) = \mathcal{U}(g) \mathcal{E}(w). \tag{II.2.6}$$

Squaring (II.2.5) and using (II.2.6) proves the first bound in (II.2.1).

For the other bound, let $w^\vee(h, \chi) = w(-h, \chi^{-1})$. Directly changing variables in the mixed Fourier coefficient gives

$$|A_{f,g}(h, \chi)| = |A_{g,f}(-h, \chi^{-1})|.$$

Consequently $\mathcal{T}_w(f, g) = \mathcal{T}_{w^\vee}(g, f)$. Inversion is an automorphism of the abelian group $G \times \widehat{G}$, so $\mathcal{E}(w^\vee) = \mathcal{E}(w)$. Apply the first bound with f and g exchanged. Equality examples establishing the final assertion are given in Section II.3. $\square$

Corollary II.2.2 (Graph version). *Let $B \subseteq G$, $\phi : B \rightarrow \widehat{G}$, and $\Gamma = \{(h, \phi(h)) : h \in B\}$. If*

$$T = \mathbb{E}_h \mathbf{1}_B(h) |\widehat{\Delta_h f}(\phi(h))|^2 > 0,$$

then, writing $E(\Gamma)$ for the unnormalized additive energy,

$$\frac{E(\Gamma)}{N^3} \geq \frac{T^4}{m(f)^4 \mathcal{U}(f)}. \quad (\text{II.2.7})$$

In particular, the hypothesis (F.7.1) implies

$$E(\Gamma) \geq \frac{\alpha^4}{m(f)^4 \mathcal{U}(f)} N^3.$$

Proof. Take $g = f$ and $w = \mathbf{1}_\Gamma$. The unnormalized Fourier coefficient is N times our normalized coefficient, so the statistic in (F.7.1), divided by N^3 , is precisely T . Positive T makes the denominator nonzero by Theorem II.2.1. $\square$

[write] **6 October 2026.** *The formal statement is recovered.* The formal `proposition_6_1` (Section `ns06_07.lean`, line 37; exact companion `proposition_6_1_holds`, `Proofs05_10.lean`, line 505) is the case $|f| \leq 1$ on $\mathbb{Z}_N$ with a single-valued graph: there $m(f) \leq 1$ and $\mathcal{U}(f) \leq m(f)^4 \leq 1$ (Lemma II.3.1), so (II.2.7) gives $E(\Gamma) \geq \alpha^4 N^3$. The intake also tested Theorem II.2.1 on the non-cyclic groups $\mathbb{Z}/2 \times \mathbb{Z}/4$ and $\mathbb{Z}/2 \times \mathbb{Z}/2 \times \mathbb{Z}/3$ and on $\mathbb{Z}/6$ (largest ratio of left to right side 0.084); the delivered suite uses cyclic groups only. *Notation:* this $\mathcal{E}(w)$ is not source 03's $\mathcal{E}(f)$ (Part III), source 04's $\mathcal{E}_j(w)$ (Section II.8) or source 05's captured energy E (Part I); see Section F.3.

II.3 The L^2 denominator and equality examples

Source 02, delivered Section 4.

Lemma II.3.1 (Two tetrahedra). *For every finite abelian group and every complex function on it,*

$$\mathcal{U}(f) \leq m(f)^4, \quad \text{equivalently} \quad \|f\|_{U^3} \leq \|f\|_2. \quad (\text{II.3.1})$$

Proof. Take the absolute value of the cube average (II.1.3). Divide its eight factors into the four vertices

$$x, \quad x + a, \quad x + b, \quad x + c$$

and the complementary four vertices

$$x + a + b, \quad x + a + c, \quad x + b + c, \quad x + a + b + c.$$

Cauchy–Schwarz bounds the cube average by the geometric mean of the averages of the squared products on these two lists. The first list is a bijective change of coordinates from G^4 to G^4 . The second is also bijective: if its coordinates are y_1, y_2, y_3, y_4 , then

$$c = y_4 - y_1, \quad b = y_4 - y_2, \quad a = y_4 - y_3, \quad x = y_1 + y_2 + y_3 - 2y_4.$$

Thus each squared-product average is $m(f)^4$, proving (II.3.1). No division by 2 occurs in the inverse map, so groups with 2-torsion are included. $\square$

Corollary II.3.2 (Variance-normalized energy). *Whenever the denominator is nonzero,*

$$\mathcal{E}(w) \geq \left(\frac{\mathcal{T}_w(f, g)}{m(f)m(g)} \right)^4.$$

For a graph and a balanced indicator $f = \mathbf{1}_A - \delta$, where $\delta = |A|/N \in (0, 1)$,

$$\boxed{\frac{E(\Gamma)}{N^3} \geq \frac{T^4}{[\delta(1-\delta)]^4 \mathcal{U}(f)} \geq \frac{T^4}{[\delta(1-\delta)]^8}.} \quad (\text{II.3.2})$$

Proof. Use Lemma II.3.1 in Theorem II.2.1. For a balanced indicator, direct averaging gives $m(f) = \delta(1-\delta)$. $\square$

The improvement in (II.3.2) is not a change in the fourth power of T . It is the retention of an exact variance denominator, together with a possibly smaller U^3 denominator. It is therefore misleading to describe it as a new exponent in Szemerédi’s theorem. It can nevertheless be a large improvement in a local energy ledger.

Example II.3.3 (Quadratic-phase graph). Let $G = \mathbb{Z}_p$ for an odd prime p , put $q(x) = \exp(2\pi i x^2/p)$, and let $f = cq$, $c \neq 0$. Then

$$\Delta_h f(x) = |c|^2 \exp(2\pi i(2hx - h^2)/p).$$

Choose the graph $h \mapsto \chi_{2h}$. We have

$$m(f) = |c|^2, \quad \mathcal{U}(f) = |c|^8, \quad T = |c|^4, \quad E(\Gamma) = N^3.$$

Equality holds in both (II.2.1) and the variance-normalized bound. In particular the coefficient 1 in the main theorem is optimal even for a single-valued graph.

Example II.3.4 (A genuinely multivalued sharp example). Let $H \leq G$ have density $d = |H|/N$, let $f = \mathbf{1}_H$, and take $w = \mathbf{1}_{H \times H^\perp}$, where $H^\perp$ is the annihilator of H . The set $H \times H^\perp$ is a subgroup of $G \times \widehat{G}$ of size N , hence $\mathcal{E}(w) = 1$. Also

$$m(f) = d, \quad \mathcal{U}(f) = d^4, \quad \mathcal{T}_w(f, f) = d^2.$$

The cube formula gives $\mathcal{U}(f) = d^4$ because all eight vertices belong to H exactly when the base and three increments do. The Fourier coefficient of $\mathbf{1}_H$ is d on $H^\perp$ and 0 elsewhere, giving the value of $\mathcal{T}$. Again equality holds in the main theorem.

Example II.3.5 (Sharpness within balanced indicators). On an even cyclic group, take A to be the even residues. Then $f = \mathbf{1}_A - 1/2$ is one half of a character. For the graph with trivial frequency at every h ,

$$m(f) = \frac{1}{4}, \quad \mathcal{U}(f) = \frac{1}{256}, \quad T = \frac{1}{16}, \quad E(\Gamma)/N^3 = 1.$$

The original T^4 bound is $1/65536$, whereas (II.3.2) gives the exact value 1. This demonstrates a strict strengthening on an admissible balanced input, not merely on arbitrary unbounded functions.

II.4 Keeping the full derivative spectrum and tracking rigidity

Source 02, delivered Section 5.

II.4.1 A support-and-energy certificate without selecting one frequency

Proposition II.4.1. *Let $f \not\equiv 0$, set $m = m(f)$ and $\kappa = \mathcal{U}(f)/m^4 \in (0, 1]$, and let $0 < \tau \leq 1$. Define*

$$\Gamma_\tau = \{(h, \chi) : |A_{f,f}(h, \chi)| \geq \tau m\}, \quad t_\tau = \frac{\mathcal{T}_{\Gamma_\tau}(f, f)}{m^2}.$$

Then

$$|\Gamma_\tau| \leq N\tau^{-2}, \tag{II.4.1}$$

$$\frac{E(\Gamma_\tau)}{N^3} \geq \frac{t_\tau^4}{\kappa}, \tag{II.4.2}$$

$$\frac{E(\Gamma_\tau)}{|\Gamma_\tau|^3} \geq \frac{t_\tau^4 \tau^6}{\kappa} \quad \text{if } \Gamma_\tau \neq \emptyset. \tag{II.4.3}$$

Proof. Parseval, followed by averaging in h , gives the exact identity

$$\sum_{h, \chi} |A_{f,f}(h, \chi)|^2 = N \mathbb{E}_{h, x} |f(x)|^2 |f(x - h)|^2 = Nm^2.$$

This proves the support bound. Apply Theorem II.2.1 to the full set Γ_τ , not a chosen graph, and substitute $\mathcal{U}(f) = \kappa m^4$. Finally divide by the cube of the support bound. Nonzero f has $\mathcal{U}(f) > 0$: the $h = 0$ summand in (II.1.2) already has a nonzero Fourier coefficient. $\square$

The certificate (II.4.3) is suitable input for an additive-energy-to-structure theorem of Balog–Szemerédi–Gowers type. The subsequent structural step is not proved anew here. The point is that there is no preliminary loss from discarding all but one frequency in each fiber. Whether that benefit survives the later graph and progression constructions is a separate quantitative problem.

II.4.2 Exact deficits in the two Cauchy–Schwarz steps

With F, D, V, K, R as in Section II.2, set

$$D_1 = m(f)^2 R - \mathcal{T}_w(f, g)^2, \quad D_2 = \mathcal{U}(g)\mathcal{E}(w) - R^2.$$

Both are nonnegative, and direct algebra gives

$$m(f)^4 \mathcal{U}(g)\mathcal{E}(w) - \mathcal{T}_w(f, g)^4 = m(f)^4 D_2 + D_1 (m(f)^2 R + \mathcal{T}_w(f, g)^2). \tag{II.4.4}$$

Thus the loss is decomposed exactly, rather than hidden in an unspecified constant.

Proposition II.4.2 (Quantitative alignment). *Suppose all quantities used to normalize below are nonzero and $0 \leq \varepsilon < 1$. If*

$$\mathcal{T}_w(f, g)^4 \geq (1 - \varepsilon) m(f)^4 \mathcal{U}(g)\mathcal{E}(w),$$

then

$$\left\| \frac{F}{\|F\|_2} - \frac{K}{\|K\|_2} \right\|_2^2 \leq 2(1 - (1 - \varepsilon)^{1/4}),$$

$$\left\| \frac{a}{\|a\|_2} - \frac{b}{\|b\|_2} \right\|_2^2 \leq 2(1 - \sqrt{1 - \varepsilon}),$$

where $a(u, \psi) = |\widehat{D}_u(\psi)|^2$ and $b(u, \psi) = |\widehat{V}_u(\psi)|^2$. The first norm uses normalized measure on G^2 ; the second uses average in u and counting measure in ψ .

Proof. Let $A = \mathcal{T}^2/(m(f)^2 R)$ and $B = R^2/(\mathcal{U}(g)\mathcal{E}(w))$. The proof of Theorem II.2.1 gives $A, B \in [0, 1]$, and the hypothesis says $A^2 B \geq 1 - \varepsilon$. Hence $A \geq \sqrt{1 - \varepsilon}$ and $B \geq 1 - \varepsilon$. The normalized inner products of the first and second pairs are $\sqrt{A}$ and $\sqrt{B}$, respectively. They are real and nonnegative, by (II.2.2) and (II.2.3). Apply $\|v - w\|^2 = 2 - 2 \operatorname{Re} \langle v, w \rangle$ for unit vectors. $\square$

[write] **6 October 2026; sources 28–33.** Near equality is now described sharply for the function: source 31's Theorem II.107.4 gives the exact optimal profile on every finite abelian group, and its Section II.109 controls the weights in the U^2 and coset topology, with examples showing that normalized L^1 control fails without complexity assumptions (item II.2).

This is an alignment theorem for the proof's intermediate functions, not a full classification of all near-extremal pairs (f, w) . The distinction matters: subgroup-localized examples and polynomial phases already occur, and broader extremizer theory is known [7].

II.4.3 An elementary exact rigidity endpoint

Proposition II.4.3. *Let $f \not\equiv 0$, and let $\phi : G \rightarrow \widehat{G}$ be defined on all of G . For its graph statistic, $T \leq m(f)^2$. Equality holds if and only if $f = cq$, with $|q| = 1$, and for each h*

$$\Delta_h q(x) = c_h \phi(h)(x), \quad |c_h| = 1.$$

In this situation ϕ is a homomorphism and all third multiplicative derivatives of q equal 1.

Proof. The chosen Fourier coefficient in each fiber has squared magnitude at most the sum of squared magnitudes in that fiber. The sum of all fibers is $m(f)^2$ after normalized averaging, by Parseval. Equality forces every $\Delta_h f$ to have Fourier support at the chosen single character. At $h = 0$, $\Delta_0 f = |f|^2$ is nonnegative and has positive mean. Its sole character must therefore be trivial, and $|f|^2 = m(f)$ everywhere. Write $f = cq$ with $|q| = 1$. Fourier inversion gives the stated derivative formula. Conversely that formula makes Parseval an equality in every fiber.

The identity

$$\Delta_{h+k} q(x) = \Delta_h q(x) \Delta_k q(x - h)$$

shows that $\phi(h+k) = \phi(h)\phi(k)$, since a nonzero scalar multiple of a character has a unique character frequency. A further derivative of a character is constant, and a derivative of a unit-modulus constant is 1; therefore third derivatives vanish multiplicatively. This definition of a quadratic phase does not impose an unjustified classical-polynomial representation in small characteristic. $\square$

II.5 Sharper common-neighborhood selection

Source 03, delivered Section 7.

Gowers' Lemma 7.4 is a convenient place to make a concrete quantitative improvement without changing the combinatorial objects in the argument. Its hypothesis controls the average intersection of a family of sets; its conclusion finds a subfamily in which almost every pair has a reasonably large intersection. The published proof samples five witnesses and estimates the useful and exceptional contributions separately [1]. Keeping those contributions together allows three witnesses to suffice.

The method remains dependent random choice, whose subsequent development is surveyed by Fox and Sudakov [11]. Our assertion here is a precise replacement for the numerical statement in the source paper. We do not assert that the resulting exponent is optimal among all possible selection arguments.

II.5.1 A cubic certificate

All pairs below are ordered, and diagonal pairs are included. For a finite family $A_1, \dots, A_n$ of subsets of a set V of cardinality $m > 0$, write

$$p_{xy} = \frac{|A_x \cap A_y|}{m}.$$

A pair will be called bad when $p_{xy} < a/2$, where $a > 0$ is the specified lower bound for the mean of the p_{xy} .

Lemma II.5.1 (Rational cubic certificate). *For every $a > 0$ and $p \geq 0$, define*

$$F_a(p) = \begin{cases} -9p^3, & 0 \leq p < a/2, \\ p^3, & p \geq a/2. \end{cases}$$

Then

$$F_a(p) \geq \frac{49}{12}a^2p - \frac{343}{108}a^3. \quad (\text{II.5.1})$$

Proof. On the good region $p \geq a/2$, the desired inequality is the tangent-line inequality for p^3 at $p = 7a/6$. More explicitly,

$$p^3 - \frac{49}{12}a^2p + \frac{343}{108}a^3 = \left(p - \frac{7a}{6}\right)^2 \left(p + \frac{7a}{3}\right) \geq 0.$$

On the interval $0 \leq p \leq a/2$, the function

$$-9p^3 - \frac{49}{12}a^2p + \frac{343}{108}a^3$$

is concave, so its minimum on that interval occurs at an endpoint. Its values at 0 and $a/2$ are, respectively,

$$\frac{343}{108}a^3 \quad \text{and} \quad \frac{-243 - 441 + 686}{216}a^3 = \frac{a^3}{108}.$$

Both are positive. At $p = a/2$, the actual function F_a takes the larger, good-region value, so the strict definition of a bad pair presents no boundary issue. $\square$

Theorem II.5.2 (Cubic replacement for Gowers' Lemma 7.4). *Let m, n be positive integers, let V have cardinality m , and let $0 < \delta \leq 1$. Suppose that $A_1, \dots, A_n \subseteq V$ satisfy*

$$\sum_{x,y=1}^n |A_x \cap A_y| \geq \delta^2 mn^2. \quad (\text{II.5.2})$$

There is a set $K \subseteq \{1, \dots, n\}$ such that

$$|K| \geq \frac{7}{3\sqrt{6}} \delta^3 n, \quad (\text{II.5.3})$$

and at least 90% of the pairs $(x, y) \in K^2$ satisfy

$$|A_x \cap A_y| \geq \frac{\delta^2 m}{2}. \quad (\text{II.5.4})$$

In particular, the conclusion holds when every A_x has cardinality at least δm .

Proof. Put $a = \delta^2$. Independently sample j_1, j_2, j_3 uniformly from V , allowing repetitions, and set

$$K = \{x : j_1, j_2, j_3 \in A_x\}.$$

Let $B(K)$ count the bad ordered pairs in K^2 . The probability that $(x, y) \in K^2$ is p_{xy}^3 . Consequently, linearity of expectation gives the exact identity

$$\begin{aligned} \mathbb{E}[|K|^2 - 10B(K)] &= \sum_{x,y=1}^n F_a(p_{xy}) \\ &\geq \frac{49}{12} a^2 \sum_{x,y=1}^n p_{xy} - \frac{343}{108} a^3 n^2 \\ &\geq \left(\frac{49}{12} - \frac{343}{108} \right) a^3 n^2 = \frac{49}{54} a^3 n^2. \end{aligned}$$

The first inequality is Lemma II.5.1; the second uses (II.5.2). Some realization therefore satisfies

$$|K|^2 - 10B(K) \geq \frac{49}{54} \delta^6 n^2 > 0.$$

It follows simultaneously that $|K| \geq \sqrt{49/54} \delta^3 n$ and $B(K) \leq |K|^2/10$. These are the required conclusions.

For the final assertion, put $d(v) = |\{x : v \in A_x\}|$. Then

$$\sum_{x,y} |A_x \cap A_y| = \sum_{v \in V} d(v)^2 \geq \frac{1}{m} \left(\sum_{v \in V} d(v) \right)^2 \geq \delta^2 mn^2,$$

by Cauchy–Schwarz and the assumed lower bounds on $|A_x|$. □

[write] **6 October 2026; sources 21–27.** Source 24's Theorem II.74.1 is the same theorem with the same certificate, proved independently; its Theorem II.75.1 states the scalar optimum below as an exact infimum.

The constant $7/(3\sqrt{6})$ is approximately 0.95258. Thus the original $2^{-1/2} \delta^5 n$ is replaced by a larger bound with exponent 3, while both the half-mean intersection threshold and the 90% requirement

are preserved. The improvement comes from a stronger estimate of the same penalized random variable, rather than from a change in the definition of success.

[write] **6 October 2026.** *Compared with the formal statement* (checked at HEAD): `lemma_7_4` (`Sections06_07.lean`, line 167) has the exact companion `lemma_7_4_holds` (`Proofs07DRC.lean`, line 536), which samples five witnesses indexed by `Fin 5`, as Section V.13 describes. Theorem II.5.2 keeps its hypothesis, the threshold $\delta^2 m/2$ and the 90% requirement, and replaces the size $2^{-1/2} \delta^5 n$ by $7\delta^3 n/(3\sqrt{6})$. Not formalized.

II.5.2 Weights and higher tuples

The certificate is independent of the counting measure on the index set. This gives a useful extension when vertices carry unequal importance or when intersections of more than two sets are under consideration.

Corollary II.5.3 (Weighted tuple selection). *Let X be a finite nonempty set, let $w_x \geq 0$ with $\sum_{x \in X} w_x = 1$, and let $(A_x)_{x \in X}$ be measurable subsets of a probability space (Ω, ν) . Fix an integer $r \geq 1$, and suppose*

$$\sum_{(x_1, \dots, x_r) \in X^r} w_{x_1} \cdots w_{x_r} \nu(A_{x_1} \cap \cdots \cap A_{x_r}) \geq a > 0.$$

There is a set $K \subseteq X$ with

$$w(K) := \sum_{x \in K} w_x \geq \left(\frac{49}{54}\right)^{1/r} a^{3/r}$$

such that the total product weight of tuples in K^r with $\nu(A_{x_1} \cap \cdots \cap A_{x_r}) < a/2$ is at most $w(K)^r/10$.

Proof. Sample three independent points of Ω and let K be the indices of the sets containing all three points. A fixed tuple lies in K^r with probability p^3 , where p is its intersection measure. Apply (II.5.1) to each tuple and sum with its product weight. The expected value of $w(K)^r$ minus ten times the bad-tuple weight is at least $49a^3/54$. Since X is finite, there are only finitely many possible sets K , and one attains at least this expectation. Taking an r th root gives the conclusion. $\square$

II.5.3 The exact scalar optimization

The simple rational point $7a/6$ was chosen to make the certificate easy to verify. One can determine the optimal supporting line for this argument. Let $z > 1$ be the unique solution of

$$2z^3 - 3z^2 = 9. \tag{II.5.5}$$

The line tangent to p^3 at $p = za/2$ passes through the limiting bad point $(a/2, -9a^3/8)$. Applying that line instead of (II.5.1) gives

$$|K| \geq \kappa_* \delta^3 n, \quad \kappa_* = \sqrt{\frac{3z^2 - z^3}{4}} = 0.9537881585 \dots \tag{II.5.6}$$

Indeed, the tangent has slope $3(za/2)^2$ and intercept $-2(za/2)^3$. Equation (II.5.5) verifies its value at the bad endpoint. Concavity on the bad interval and convexity on the good interval again prove a global minorant. Its value at the mean a is $(3z^2 - z^3)a^3/4$. Existence and uniqueness of $z > 1$ follow from the fact that $2z^3 - 3z^2$ is strictly increasing on $(1, \infty)$, starts at -1 , and tends to infinity. Moreover $2 < z < 5/2$, so the expression under the square root is positive.

For clarity, the meaning of optimality here is limited and exact. When $za/2 \leq 1$, consider arbitrary scalar random variables $P \in [0, 1]$ with $\mathbb{E}P = a$. A two-point distribution, with one value tending to $(a/2)^-$ and the other equal to $za/2$, approaches the tangent-line lower bound for $\mathbb{E}F_a(P)$. Thus no larger bound follows from this scalar mean constraint alone. The infimum is approached because badness is defined by a strict inequality.

This observation also explains a restriction of a two-witness version of the same argument. Its penalized scalar function is p^2 on the good region and $-9p^2$ on the bad region. The analogous tangent point is ha , where $h = (1 + \sqrt{10})/2$. The tangent's value at a is

$$(2h - h^2)a^2 < 0.$$

For sufficiently small a , the corresponding two-point distribution lies in $[0, 1]$ and gives a negative limiting expectation. Consequently, two witnesses and this scalar first-moment certificate cannot produce a positive uniform lower bound at the original parameters. This does not exclude a better theorem exploiting properties of actual intersection matrices, nor does it prove optimality of the exponent 3.

II.5.4 A flexible one-witness alternative

For applications, keeping the threshold exactly at half the mean may be unnecessary. Allowing a smaller constant multiple of the mean gives a much larger selected set using just one witness. This is an explicit form of the established variant in Fox and Sudakov [11, Lemma 5.1].

Lemma II.5.4 (One-witness tradeoff). *Under (II.5.2), let $0 < \gamma < \varepsilon < 1$. There is a set $K \subseteq \{1, \dots, n\}$ satisfying*

$$|K| \geq \sqrt{1 - \gamma/\varepsilon} \delta n$$

such that at most an ε fraction of the ordered pairs in K^2 have intersection smaller than $\gamma\delta^2 m$.

Proof. Choose $j \in V$ uniformly and set $K = \{x : j \in A_x\}$. Then

$$\mathbb{E}|K|^2 = \sum_{x,y} p_{xy} \geq \delta^2 n^2.$$

If $B(K)$ counts the pairs in K^2 with $p_{xy} < \gamma\delta^2$, then

$$\mathbb{E}B(K) = \sum_{p_{xy} < \gamma\delta^2} p_{xy} \leq \gamma\delta^2 n^2.$$

Thus

$$\mathbb{E}[|K|^2 - \varepsilon^{-1}B(K)] \geq (1 - \gamma/\varepsilon)\delta^2 n^2 > 0.$$

A realization attaining at least this expectation gives both claims. □

II.6 An elementary quantitative replacement for Proposition 7.3

Source 03, delivered Section 8.

The preceding selection statements have different uses. The cubic version strengthens the exact statement of Lemma 7.4. The one-witness version is more efficient for the later small-difference-set

argument, because that argument can tolerate a smaller absolute constant in the common-neighbor threshold. We also remove a separate loss by counting the total mass of popular differences directly.

Throughout this section, A is a finite nonempty subset of an arbitrary abelian group and $m = |A|$. Put

$$r_A(d) = \left| \{(a, b) \in A^2 : a - b = d\} \right|, \quad E(A) = \sum_d r_A(d)^2.$$

We assume $E(A) \geq cm^3$, where $0 < c \leq 1$. The sums involve only finitely many nonzero terms, so no finiteness assumption on the ambient group is required.

II.6.1 Popular differences retain more edge mass

Lemma II.6.1 (Popular mass). *Let*

$$D = \{d : r_A(d) \geq cm/2\}, \quad M = \sum_{d \in D} r_A(d), \quad \rho = \frac{c}{2-c}.$$

Then $D = -D$, $|D| \leq 2m/c$, and $M \geq \rho m^2$.

Proof. The symmetry follows from $r_A(d) = r_A(-d)$. Since $\sum_d r_A(d) = m^2$, the lower bound $r_A(d) \geq cm/2$ on D gives $|D| \leq 2m/c$. On D we have $r_A(d)^2 \leq mr_A(d)$, while off D we have $r_A(d)^2 \leq (cm/2)r_A(d)$. Therefore

$$cm^3 \leq E(A) \leq mM + \frac{cm}{2}(m^2 - M).$$

Rearranging proves $M \geq cm^2/(2-c)$. □

This estimate retains the edge density of the popular-difference relation. It avoids the loss from first counting popular values and then replacing each of their multiplicities by the smallest allowed multiplicity. In particular, the relation has density at least $\rho \geq c/2$.

II.6.2 A general walk-counting statement

Proposition II.6.2 (Compression by four edge differences). *Let $G \subseteq A \times A$ be a symmetric relation, with loops allowed, satisfying $|G| \geq \rho m^2$, where $0 < \rho \leq 1$. Suppose that every $(a, b) \in G$ has $a - b \in D$ for a finite set D , and put $L = |D|$. For $0 < \gamma < \varepsilon < \beta < 1/2$, set $s = \sqrt{1 - \gamma/\varepsilon}$. There is $A' \subseteq A$ such that*

$$|A'| \geq (1 - \varepsilon/\beta)s\rho m, \tag{II.6.1}$$

$$|A' - A'| \leq \frac{L^4}{(1 - 2\beta)\gamma^2 s \rho^5 m^3}. \tag{II.6.2}$$

Proof. Write $N(a) = \{v : (a, v) \in G\}$. By symmetry, double-counting, and Cauchy–Schwarz,

$$\sum_{a, b \in A} |N(a) \cap N(b)| = \sum_{v \in A} |N(v)|^2 \geq \frac{|G|^2}{m} \geq \rho^2 m^3.$$

Lemma II.5.4 supplies $U \subseteq A$, of cardinality $u \geq s\rho m$, with at most εu^2 bad ordered pairs. Here a pair is bad when its common G -neighborhood has cardinality smaller than $\gamma\rho^2 m$.

Let A' consist of the vertices of U that belong, in the first coordinate, to at most βu bad pairs in U^2 . Counting those pairs shows that at most $(\varepsilon/\beta)u$ vertices are excluded. This proves (II.6.1). Each $a \in A'$ has at least $(1 - \beta)u$ good partners in U . Consequently, every $a, b \in A'$ have at least $(1 - 2\beta)u$ common good partners $k \in U$.

Fix such a, b . For every common good partner k , choose

$$x \in N(a) \cap N(k), \quad y \in N(b) \cap N(k).$$

There are at least $(1 - 2\beta)u(\gamma\rho^2m)^2$ triples (k, x, y) . Associate to each the four edge differences

$$d_1 = a - x, \quad d_2 = k - x, \quad d_3 = k - y, \quad d_4 = b - y.$$

They lie in D and satisfy

$$a - b = d_1 - d_2 + d_3 - d_4.$$

For the fixed endpoints a, b , the association is injective: one recovers successively $x = a - d_1$, $k = x + d_2$, and $y = k - d_3$.

For each element of $A' - A'$, fix one pair of endpoints representing it. The associated label quadruples for distinct differences are disjoint, because a label quadruple determines $d_1 - d_2 + d_3 - d_4$. Since there are only L^4 possible label quadruples,

$$|A' - A'| (1 - 2\beta)u\gamma^2\rho^4m^2 \leq L^4.$$

Using $u \geq spm$ proves (II.6.2). This counts walks with specified endpoints; repeated vertices and loops are allowed throughout, and cause no loss in the injection. $\square$

II.6.3 Explicit bounds and their scope

Theorem II.6.3 (Elementary historical improvement). *If $E(A) \geq c|A|^3$ with $0 < c \leq 1$, there is $A' \subseteq A$ satisfying*

$$|A'| \geq \frac{c}{8}|A|, \quad |A' - A'| \leq 2^{19}c^{-9}|A|. \quad (\text{II.6.3})$$

More precisely, with $\rho = c/(2 - c)$, one may take

$$|A'| \geq \frac{\rho m}{2\sqrt{2}}, \quad |A' - A'| \leq 6912\sqrt{2}\rho^{-5}c^{-4}m. \quad (\text{II.6.4})$$

Proof. Use D from Lemma II.6.1, and define $G = \{(a, b) \in A^2 : a - b \in D\}$. That lemma gives $|G| \geq \rho m^2$ and $L \leq 2m/c$. Apply Proposition II.6.2 with

$$\gamma = \frac{1}{12}, \quad \varepsilon = \frac{1}{6}, \quad \beta = \frac{1}{3}, \quad s = \frac{1}{\sqrt{2}}.$$

The size bound is $\rho m/(2\sqrt{2})$, and the difference bound becomes

$$|A' - A'| \leq \frac{432\sqrt{2}L^4}{\rho^5m^3} \leq 6912\sqrt{2}\rho^{-5}c^{-4}m.$$

Finally $\rho \geq c/2$, so

$$|A'| \geq \frac{cm}{4\sqrt{2}} \geq \frac{cm}{8}, \quad |A' - A'| \leq 221184\sqrt{2}c^{-9}m < 2^{19}c^{-9}m.$$

This proves both versions. $\square$

[write] **6 October 2026; sources 21–27.** Source 24’s Corollary II.77.2 gives $(\kappa/8, 2^{15}\kappa^{-5})$ with relative doubling $2^{17}\kappa^{-6}$; source 10’s Theorem II.20.1 remains the best of the three.

For comparison, Proposition 7.3 in the source paper gives a subset of size at least $2^{-20}c^{12}m$ with difference set of size at most $2^{38}c^{-24}m$ [1]. Theorem II.6.3 improves both exponents using only one sampled vertex, Cauchy–Schwarz, and the explicit counting injection above. Even if one changes only Lemma 7.4 to Theorem II.5.2 and leaves the remaining original proof intact, its formulas yield the weaker but direct improvements

$$|A'| \geq 2^{-14}c^8m, \quad |A' - A'| \leq 2^{31}c^{-20}m.$$

Indeed, the original parameter $\delta = c^2/8$ then gives selected density $\alpha = \kappa\delta^4$, where $\kappa = 7/(3\sqrt{6})$; the original cleanup and walk count give $\alpha m/2$ and $120m/(\alpha\delta^4c^4)$, respectively. The displayed relaxations use $\kappa > 1/2$ and $120/\kappa < 128$.

[write] **6 October 2026; sources 06–10.** *A stronger elementary replacement (source 10).* Under the same hypothesis $E(A) \geq c|A|^3$, $0 < c \leq 1$, source 10’s Theorem II.20.1 gives $|A_*| \geq cm/4$ and $|A_* - A_*| \leq 2^{10}c^{-5}m$, on any abelian group, by weighted dependent random choice with the diagonal pairs kept. It is stronger than (II.6.3) in both conclusions: $c/4 > c/8$, and $2^{10}c^{-5} \leq 2^{19}c^{-9}$ because $c^4 \leq 1 \leq 2^9$. Both are far weaker than Reiher–Schoen [12], as both sources say. Both imply the formal `proposition_7_3` bounds $2^{-20}c^{12}$ and $2^{38}c^{-24}$ (for source 10: $c^{11} \leq 2^{18}$ and $c^{19} \leq 2^{28}$); neither is formalized. Lemma 7.4 is treated only here.

These are improvements to the historical proof architecture, rather than to the strongest available Balog–Szemerédi–Gowers theorem. For example, Reiher and Schoen prove that, for every $0 < \eta < 1/2$, the same energy hypothesis gives a subset satisfying

$$|A'| \geq (1 - \eta)c^{1/2}m, \quad |A' - A'| \leq 2^{33}\eta^{-9}c^{-4}|A'|$$

[12, Theorem 1.2]. Their result is substantially stronger. The purpose of the elementary version is to supply a short, fully explicit replacement that can be checked independently and integrated into the existing development with a small dependency set. A complete improvement to the final quantitative Szemerédi bound would also require propagating the new parameters through the later arguments; no such consequence follows from this section alone.

[write] **6 October 2026.** *Compared with the formal statement (checked at HEAD):* `proposition_7_3` (Sections06_07.lean, line 135; exact companion `proposition_7_3_holds`, Proofs07BalogSzemeredi.1 ean, line 996) is stated on $\mathbb{Z}^n$ with the bounds $2^{-20}c^{12}|A|$ and $2^{38}c^{-24}|A|$. For $0 < c \leq 1$, $c/8 \geq 2^{-20}c^{12}$ (as $c^{11} \leq 2^{17}$) and $2^{19}c^{-9} \leq 2^{38}c^{-24}$ (as $c^{15} \leq 2^{19}$), so Theorem II.6.3, which holds in every abelian group, implies the bounds of the formal statement. Not formalized.

II.7 Notation and normalization

Source 04, delivered Section 2.

Let G be a finite abelian group, written additively, and let $N = |G|$. We write $\mathbb{E}_{x \in G} g(x) = N^{-1} \sum_{x \in G} g(x)$. Characters are written multiplicatively. Except in the moment section, which explicitly introduces an unnormalized transform, Fourier coefficients are normalized:

$$\widehat{f}(\chi) = \mathbb{E}_x f(x) \overline{\chi(x)}.$$

Thus

$$\mathbb{E}_x |f(x)|^2 = \sum_{\chi \in \widehat{G}} |\widehat{f}(\chi)|^2.$$

Sums over the dual group in this convention are ordinary sums, not averages.

For $d \geq 1$, put

$$\|f\|_{U^d(G)}^{2^d} = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{\epsilon \in \{0,1\}^d} \mathcal{C}^{|\epsilon|} f(x + \epsilon_1 h_1 + \dots + \epsilon_d h_d),$$

where $\mathcal{C}$ is complex conjugation. We omit G when it is clear. For $d = 1$ this means $\|f\|_{U^1} = |\mathbb{E}f|$. For $d \geq 2$ it is the Gowers norm. If f is real, conjugations can be omitted.

For a real function w , define its normalized cube count

$$Q_d(w) = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{\epsilon \in \{0,1\}^d} w(x + \epsilon \cdot h).$$

For $w = 1_A$, $N^{d+1}Q_d(w)$ counts all ordered additive d -cubes, including degenerate cubes, exactly as in the corresponding count in the source. If $\delta = |A|/N$, its balanced function is $f = 1_A - \delta$.

The conversion to Gowers's parameter is important:

$$f \text{ is } \alpha\text{-uniform of degree } d-1 \iff \|f\|_{U^d}^{2^d} \leq \alpha.$$

In particular, the letter u below generally denotes a norm, while α denotes an upper bound on a power of that norm. Confusing these two quantities would conceal the improvement in the exponent.

We use the standard Gowers–Cauchy–Schwarz inequality in normalized form:

$$\left| \mathbb{E}_{x, h} \prod_{\epsilon} \mathcal{C}^{|\epsilon|} f_{\epsilon}(x + \epsilon \cdot h) \right| \leq \prod_{\epsilon} \|f_{\epsilon}\|_{U^d}.$$

It follows by applying Cauchy–Schwarz successively in the cube coordinates; this is Lemma 3.8 in the source. Its ordinary consequences include the triangle inequality and $\|f\|_{U^d} \leq \|f\|_{U^{d+1}}$. For clarity, the latter also follows by writing

$$\|f\|_{U^d}^{2^d} = \mathbb{E}_{h_1, \dots, h_{d-1}} \left| \mathbb{E}_x \Delta_{h_1} \cdots \Delta_{h_{d-1}} f(x) \right|^2.$$

Square this identity and use Cauchy–Schwarz, then apply $|\mathbb{E}g|^4 \leq \|g\|_{U^2}^4$ to each derivative. The resulting right-hand side is $\|f\|_{U^{d+1}}^{2^{d+1}}$.

An ordered additive $2d$ -tuple of a set B is a tuple with $x_1 + \dots + x_d = x_{d+1} + \dots + x_{2d}$. A map $\phi : B \rightarrow H$ *respects* the tuple if the same equation holds after applying ϕ . All tuple counts in this article are ordered and allow repetitions unless expressly called nonexceptional.

II.8 Exact second moments and density-sensitive lifting

Source 04, delivered Section 3.

Several losses in the paper occur when a second moment is bounded using the size of the ambient space rather than the size of the set under consideration. Retaining the exact second moment is useful on its own, and becomes particularly effective when followed by the restriction argument in the next section. The underlying interpolation inequality is standard Hölder; the contribution here is the careful normalization, the resulting sharper statements, and their use in the subsequent selection theorem.

II.8.1 Weighted energies on finite abelian groups

Let H be a finite abelian group of order M , and let $\widehat{H}$ be its character group. For $w : H \rightarrow \mathbb{C}$, use the unnormalized transform

$$\widehat{w}(\chi) = \sum_{x \in H} w(x) \overline{\chi(x)}.$$

For an integer $j \geq 1$, define

$$\mathcal{E}_j(w) = \frac{1}{M} \sum_{\chi \in \widehat{H}} |\widehat{w}(\chi)|^{2j}.$$

Character orthogonality gives

$$\mathcal{E}_1(w) = \sum_{x \in H} |w(x)|^2.$$

For an indicator $w = 1_A$, the same orthogonality calculation shows that $\mathcal{E}_j(A)$ counts ordered tuples $(a_1, \dots, a_{2j}) \in A^{2j}$ satisfying

$$a_1 + \dots + a_j = a_{j+1} + \dots + a_{2j}.$$

Repeated entries are included throughout.

Theorem II.8.1 (Exact energy interpolation). *Let $j \geq 2$, and suppose that $\mathcal{E}_1(w) > 0$. Then*

$$\boxed{\mathcal{E}_j(w) \geq \frac{\mathcal{E}_2(w)^{j-1}}{\mathcal{E}_1(w)^{j-2}}.} \quad (\text{M1})$$

For $j > 2$, equality holds if and only if all nonzero Fourier coefficients of w have the same modulus. In particular, if $A \neq \emptyset$, then

$$\mathcal{E}_j(A) \geq \frac{\mathcal{E}_2(A)^{j-1}}{|A|^{j-2}}. \quad (\text{M2})$$

For $j > 2$, equality in (M2) holds exactly when A is a coset of a subgroup.

Proof. Put $x_\chi = |\widehat{w}(\chi)|^2$. If $j = 2$, the assertion is an identity. For $j > 2$, Hölder gives

$$\sum_{\chi} x_{\chi}^2 = \sum_{\chi} x_{\chi}^{(j-2)/(j-1)} x_{\chi}^{j/(j-1)} \leq \left(\sum_{\chi} x_{\chi} \right)^{(j-2)/(j-1)} \left(\sum_{\chi} x_{\chi}^j \right)^{1/(j-1)}.$$

Raising to the power $j-1$ and dividing by M^{j-1} proves (M1). Equality in Hölder says that the positive values of x_{χ} are equal.

Now let $w = 1_A$, and put $m = |A|$. Its trivial Fourier coefficient is m . Thus equality forces every nonzero Fourier coefficient to have modulus m . Let K be the subgroup generated by $A - A$. Equality in the triangle inequality shows that $|\widehat{1_A}(\chi)| = m$ precisely when χ is constant on A , equivalently when $\chi \in K^{\perp}$. Consequently Parseval gives

$$Mm = |K^{\perp}| m^2 = \frac{M}{|K|} m^2.$$

Hence $m = |K|$. Since A is contained in a coset of K , it is that entire coset. Conversely, the Fourier transform of a coset indicator has constant modulus on its support, so it gives equality. $\square$

[write] 6 October 2026; sources 15–20. Source 17’s Proposition II.59.1 and Corollary II.59.2 prove (M2) with its coset equality case and (M3) again, independently; its remark on the exact second moment $N^2|B|$ is (M4)–(M5). (M6) has no counterpart there.

The proof applies without change to any finite family of spectra: the index χ can be replaced by a pair (i, χ) . This observation will be used for arrangements, where the extra index records common sidelengths.

Corollary II.8.2 (Lifting respected graph energies). *Let B be a nonempty subset of a finite abelian group, and let $\phi : B \rightarrow K$ take values in another finite abelian group. Write $m = |B|$, let $0 \leq \gamma \leq 1$, and suppose that at least γm^3 ordered additive quadruples of B are respected by ϕ . Then, for every integer $j \geq 2$, at least*

$$\gamma^{j-1} m^{2j-1} \tag{M3}$$

ordered additive $2j$ -tuples are respected by ϕ .

Proof. Apply Theorem II.8.1 to the indicator of the graph $\Gamma = \{(x, \phi(x)) : x \in B\}$. Its first energy is m , and its j -th energy is exactly the number in question. $\square$

For $B \subset \mathbb{F}_p$ of size βp , the source proof of Lemma 9.2 uses

$$\sum_{u,r} \left| \widehat{f}_u(r) \right|^2 \leq p^3.$$

The exact identity is

$$\sum_{u,r} \left| \widehat{f}_u(r) \right|^2 = p^2 |B| = \beta p^3. \tag{M4}$$

Thus, under its unnormalized hypothesis of at least $\gamma_0 p^3$ respected quadruples, its conclusion strengthens to $\gamma_0^7 \beta^{-6} p^{15}$ respected 16-tuples. Under the hypothesis of Corollary 9.4, where $\gamma_0 = \gamma \beta^3$, this is

$$\gamma^7 \beta^{15} p^{15}, \tag{M5}$$

rather than $\gamma^7 \beta^{21} p^{15}$. These comparisons refer to the displayed statements in the source and its repository transcription [1, 23].

[write] 6 October 2026. *Proof-level baseline* (checked at HEAD): in `Proofs09Moments.lean`, the private lemma `section9_second_moment_le` (line 480), used for `lemma_9_2_holds` (line 527), computes the exact value $N^2|B|$ of this sum (line 495) before weakening it to N^3 . So (M4) is already a proof-level fact in the repository, as the stronger Corollary 2.5 pair is (Part I); the exported `lemma_9_2` keeps the weaker bound, and (M5) and Corollary II.8.2 are not formalized.

II.8.2 A finite-size correction from the trivial character

There is a further refinement when the ambient group size matters. Its proof also identifies a useful limitation of using only the full first and second moments.

Proposition II.8.3 (Removing the trivial Fourier coefficient). *Let $A \subset H$, where $0 < m = |A| < M = |H|$, and put*

$$\rho = \frac{m}{M}, \quad \gamma = \frac{\mathcal{E}_2(A)}{m^3}.$$

Then $\gamma \geq \rho$, and for every integer $j \geq 2$,

$$\frac{\mathcal{E}_j(A)}{m^{2j-1}} \geq \rho + \frac{(\gamma - \rho)^{j-1}}{(1 - \rho)^{j-2}}. \quad (\text{M6})$$

The right side is at least γ^{j-1} . If only $\mathcal{E}_2(A) \geq \gamma_0 m^3$ is known, one may replace $\gamma - \rho$ by $\max\{\gamma_0 - \rho, 0\}$.

Proof. Remove the trivial character before applying the Hölder argument. The sums of the remaining squared and fourth Fourier moduli are

$$M\left(m - \frac{m^2}{M}\right), \quad M\left(\mathcal{E}_2(A) - \frac{m^4}{M}\right),$$

respectively. The remaining $2j$ -th moment is $M(\mathcal{E}_j(A) - m^{2j}/M)$. Substitution in the same interpolation inequality proves (M6). The second displayed quantity is nonnegative, giving $\gamma \geq \rho$. Finally, write

$$\gamma = \rho \cdot 1 + (1 - \rho) \frac{\gamma - \rho}{1 - \rho}$$

and apply convexity of $t \mapsto t^{j-1}$. □

This correction can be sharp away from subgroup examples. For the parabola $\Gamma = \{(x, x^2) : x \in \mathbb{F}_p\}$, with p odd, the trivial Fourier coefficient has modulus p , the coefficients at $(r, 0)$ with $r \neq 0$ vanish, and all coefficients whose second frequency is nonzero have modulus $\sqrt{p}$. Indeed, completing the square reduces the latter assertion to

$$\left| \sum_x e_p(ux^2) \right|^2 = \sum_{x,y} e_p(u(x-y)(x+y)) = p \quad (u \neq 0),$$

where $e_p(t) = \exp(2\pi it/p)$; the change of variables $(x, y) \mapsto (x - y, x + y)$ is invertible. It follows that

$$\mathcal{E}_2(\Gamma) = 2p^2 - p, \quad \mathcal{E}_j(\Gamma) = p^{2j-2} + (p-1)p^{j-1}. \quad (\text{M7})$$

All nonzero residual squared Fourier moduli agree, so Proposition II.8.3 is an equality in this example.

II.9 Polynomial restriction to approximate homomorphisms

Source 04, delivered Section 4.

We now replace the repeated degree-one filters in the proof of Lemma 9.3 by one filter of higher degree. The degree is chosen from the desired accuracy. Although more coefficient relations appear, their nonnegative Fourier weights have total mass one. A weighted error bound therefore avoids a loss proportional to the number of those relations. The retained number of configurations has a polynomial dependence with exponent linear in the number of entries, instead of the much larger exponent in the source.

For $B \subset \mathbb{F}_p$, let $T_d(B)$ denote its number of ordered additive $2d$ -tuples, and let $T_d^\phi(B)$ denote the number also satisfying

$$\phi(x_1) + \cdots + \phi(x_d) = \phi(x_{d+1}) + \cdots + \phi(x_{2d}). \quad (\text{R1})$$

Thus $T_d^\phi(B) = \mathcal{E}_d(\Gamma)$. An elementary upper bound that will be used repeatedly is

$$T_d(B) \leq |B|^{2d-1}, \quad (\text{R2})$$

because the first $2d - 1$ entries uniquely determine the last.

Theorem II.9.1 (Polynomial restriction theorem). *Let $d \geq 2$ be an integer, and let $0 < \alpha, \beta, \eta \leq 1$. Define*

$$\kappa_d = \frac{1}{2(400d)^d}, \quad \Lambda_d = 2 \left(1 + \binom{2d}{2} \right) (400d)^d. \quad (\text{R3})$$

Suppose that p is a prime with

$$p \geq \Lambda_d(\alpha\eta)^{-2d}\beta^{-1}. \quad (\text{R4})$$

Let $B \subset \mathbb{F}_p$, with $|B| = \beta p$, and let $\phi : B \rightarrow \mathbb{F}_p$ satisfy

$$T_d^\phi(B) \geq \alpha\beta^{2d-1}p^{2d-1}. \quad (\text{R5})$$

Then there is $B' \subset B$ for which

$$T_d(B') \geq \kappa_d \alpha^{2d} \eta^{2d-1} \beta^{2d-1} p^{2d-1}, \quad (\text{R6})$$

and

$$\frac{T_d^\phi(B')}{T_d(B')} \geq \frac{1}{1 + \eta} \geq 1 - \eta. \quad (\text{R7})$$

[write] **6 October 2026; sources 15–20.** Source 17's Theorem II.61.2 is a complementary prime-field restriction theorem with the same powers $\alpha^{2d}\eta^{2d-1}$: a better constant for $d \leq 97$ ($1/(2(4(d+1)^2)^d)$ against κ_d), a worse size threshold (degree $4d - 1$ in $(\alpha\eta)^{-1}$ and $\beta^{-(2d-1)}$, against (R4)); comparison in the note after that theorem. Its fixed-characteristic Section II.63 has no counterpart here.

[write] **6 October 2026; sources 28–33.** Over a fixed field source 29's Theorem II.120.4 lowers the threshold to $8qS/(\alpha\eta)$, linear in $(\alpha\eta)^{-1}$, keeping every degenerate tuple, and its Corollary II.120.5 removes it; its constants depend on the field (item II.6).

Equivalently, the size hypothesis is $|B| \geq \Lambda_d(\alpha\eta)^{-2d}$. It depends on the number of available points rather than a separate density requirement. The theorem concerns preservation of additive configurations. It does not merely assert that B' has many points. The normalization by $|B|^{2d-1}$, rather than the ambient configuration count, is essential for its use after Corollary II.8.2.

II.9.1 The kernel and its coefficients

Lemma II.9.2 (A powered cosine with many useful coefficients). *Let d, L be positive integers, and put $q = 4dL^2$. Write*

$$K_q(t) = \cos^{2q}(\pi t) = \sum_{\ell=-q}^q b_\ell e^{2\pi i \ell t}, \quad b_\ell = 4^{-q} \binom{2q}{q+\ell}. \quad (\text{R8})$$

Then $0 \leq K_q \leq 1$, its coefficients are nonnegative and symmetric, and

$$b_0 \geq \frac{1}{2\sqrt{q}}, \quad \sum_{\ell=-q}^q \left(\frac{b_\ell}{b_0} \right)^{2d} \geq L + 1. \quad (\text{R9})$$

Proof. The expansion follows by applying the binomial theorem to $(e^{\pi it} + e^{-\pi it})^{2q}/2^{2q}$. For the central coefficient, write $u_q = 4^{-q} \binom{2q}{q}$. We have $u_1 = 1/2$ and

$$\frac{u_{q+1}}{u_q} = \frac{2q+1}{2q+2} \geq \sqrt{\frac{q}{q+1}}.$$

After squaring, the last inequality reduces to $(2q+1)^2 \geq 4q(q+1)$. Induction proves the stated lower bound for b_0 .

For $0 \leq \ell \leq L$,

$$\frac{b_\ell}{b_0} = \prod_{j=1}^{\ell} \frac{q-j+1}{q+j} = \prod_{j=1}^{\ell} \left(1 - \frac{2j-1}{q+j}\right) \geq 1 - \frac{\ell^2}{q}. \quad (\text{R10})$$

Here all factors lie in $[0, 1]$, and $\prod_j (1 - z_j) \geq 1 - \sum_j z_j$, while $\sum_{j=1}^{\ell} (2j-1)/(q+j) \leq \ell^2/q$. Bernoulli's inequality now gives

$$\left(\frac{b_\ell}{b_0}\right)^{2d} \geq 1 - \frac{2d\ell^2}{q} \geq \frac{1}{2}. \quad (\text{R11})$$

The central term contributes 1, and the $2L$ terms with $1 \leq |\ell| \leq L$ contribute at least L . $\square$

II.9.2 Repeated entries and weighted relation errors

A direct union bound over all the kernel's coefficient relations would be expensive, because its degree depends on the desired accuracy. Instead, we keep their Fourier weights. Positivity and the fact that the weights sum to one allow all the unwanted relations to be controlled together.

Lemma II.9.3 (A rank-two relation inside a finite set). *Let $d \geq 2$, let p be prime, let $B \subset \mathbb{F}_p$ have size $m > 0$, and put*

$$v = (\underbrace{1, \dots, 1}_d, \underbrace{-1, \dots, -1}_d) \in \mathbb{F}_p^{2d}.$$

For every $\epsilon \in \mathbb{F}_p^{2d} \setminus \mathbb{F}_p v$,

$$\#\{x \in B^{2d} : v \cdot x = \epsilon \cdot x = 0\} \leq m^{2d-2}. \quad (\text{R12})$$

The number of additive $2d$ -tuples in B with repeated entries is at most $\binom{2d}{2} m^{2d-2}$.

Proof. The matrix with rows v, ϵ has rank two, so it has an invertible 2-by-2 column minor. Fix the other $2d-2$ coordinates arbitrarily in B . The two remaining coordinates are uniquely determined in $\mathbb{F}_p$; either both belong to B , or this choice gives no admissible completion. This proves (R12).

If $x_i = x_j$, the vector $e_i - e_j$ supplies a second relation. It is not a scalar multiple of v , since it has zero coordinates whereas every nonzero scalar multiple of v has none. Apply (R12) to each of the $\binom{2d}{2}$ pairs of positions and take a union bound. $\square$

Lemma II.9.4 (Weighted selection estimates). *Let $d \geq 2$, $q \geq 1$, and let $p > 2q$ be prime. Let $B \subset \mathbb{F}_p$ have size $m > 0$, and let $\phi : B \rightarrow \mathbb{F}_p$. Choose r, s independently and uniformly in $\mathbb{F}_p$. Conditional on this choice, retain each point $x \in B$ independently with probability*

$$K_q \left(\frac{rx + s\phi(x)}{p} \right). \quad (\text{R13})$$

Integer representatives do not matter because K_q is periodic. Put

$$P_0 = b_0^{2d}, \quad Q = \sum_{\ell=-q}^q \left(\frac{b_\ell}{b_0} \right)^{2d}, \quad C_d = 1 + \binom{2d}{2}. \quad (\text{R14})$$

Let X count all selected respected additive tuples, and let Y count all selected unrespected additive tuples. Then

$$\mathbb{E}X \geq P_0 Q T_d^\phi(B), \quad \mathbb{E}Y \leq P_0 (T_d(B) - T_d^\phi(B)) + C_d m^{2d-2}. \quad (\text{R15})$$

Proof. For a fixed tuple $x = (x_1, \dots, x_{2d})$, abbreviate its point probabilities by $P_{r,s}(x_j)$, and form the product with multiplicities

$$A_{r,s}(x) = \prod_{j=1}^{2d} P_{r,s}(x_j).$$

Its actual conditional survival probability is the product over the *distinct* entries of x . Since all point probabilities lie in $[0, 1]$, the actual probability is at least $A_{r,s}(x)$, equals it when the entries are distinct, and exceeds it by at most 1 otherwise. This observation accounts for repeated points without assuming that their inclusion events are independent.

Expand $A_{r,s}(x)$ using the nonnegative Fourier coefficients from (R8). For $\epsilon \in \{-q, \dots, q\}^{2d}$, put $w_\epsilon = \prod_j b_{\epsilon_j}$. Orthogonality gives

$$\mathbb{E}_{r,s} A_{r,s}(x) = \sum_{\epsilon} w_\epsilon 1_{\{\epsilon \cdot x = 0\}} 1_{\{\epsilon \cdot \phi(x) = 0\}}, \quad \sum_{\epsilon} w_\epsilon = \left(\sum_{\ell=-q}^q b_\ell \right)^{2d} = 1. \quad (\text{R16})$$

Here $\epsilon \cdot \phi(x)$ denotes $\sum_j \epsilon_j \phi(x_j)$. The last equality uses $K_q(0) = 1$. All weights in this identity are nonnegative.

If ϵ is a scalar multiple of v modulo p , the condition $p > 2q$ forces its integer entries to be $(\ell, \dots, \ell, -\ell, \dots, -\ell)$ for a single $\ell \in [-q, q]$. These terms have total weight $P_0 Q$. Every one contributes for a respected additive tuple. Thus its expected product with multiplicities is at least $P_0 Q$, and the same lower bound holds for its actual survival probability. Summing proves the assertion about $\mathbb{E}X$.

For an unrespected additive tuple, every nonzero such ℓ is invertible modulo p , so only $\ell = 0$ can contribute among these scalar multiples. Its weight is P_0 . For each remaining coefficient vector, Lemma II.9.3 bounds the number of domain tuples satisfying the additional relation by m^{2d-2} . Requiring that they be unrespected, or imposing the additional image relation in (R16), can only reduce this number. Summing first over tuples and then over coefficient vectors, and using $\sum_{\epsilon} w_\epsilon = 1$, gives

$$\sum_{\substack{x \in B^{2d} \\ x \text{ additive and unrespected}}} \mathbb{E}_{r,s} A_{r,s}(x) \leq P_0 (T_d(B) - T_d^\phi(B)) + m^{2d-2}.$$

Finally, the difference between actual survival and this product need only be charged to repeated tuples. Lemma II.9.3 bounds their number by $\binom{2d}{2} m^{2d-2}$. Adding that error proves the assertion about $\mathbb{E}Y$. $\square$

II.9.3 Proof of the polynomial restriction theorem

Proof of Theorem II.9.1. Set

$$a = \alpha\eta, \quad m = |B| = \beta p, \quad H = m^{2d-1}, \quad c = \kappa_d a^{2d}. \quad (\text{R17})$$

Choose

$$L = \left\lceil \frac{4}{a} \right\rceil, \quad q = 4dL^2. \quad (\text{R18})$$

Since $a \leq 1$,

$$L \leq \frac{5}{a}, \quad q \leq \frac{100d}{a^2}. \quad (\text{R19})$$

The size condition is $m \geq \Lambda_d a^{-2d}$. It implies $p > 2q$, since $p \geq m$, $\Lambda_d > 200d$, and $a^{-2d} \geq a^{-2}$. Furthermore, with $C_d = 1 + \binom{2d}{2}$, it gives

$$C_d m^{2d-2} \leq cH, \quad (\text{R20})$$

because $\Lambda_d = C_d/\kappa_d$.

Apply Lemma II.9.4 with this kernel. Lemma II.9.2 gives

$$Q \geq L + 1 \geq \frac{4}{a}, \quad P_0 \geq (4q)^{-d} \geq \frac{a^{2d}}{(400d)^d} = 2c. \quad (\text{R21})$$

The input hypothesis and the bound $T_d(B) \leq H$ therefore imply

$$\mathbb{E}X \geq P_0 Q \alpha H, \quad \mathbb{E}Y \leq P_0 H + C_d m^{2d-2}. \quad (\text{R22})$$

Consequently

$$\begin{aligned} \mathbb{E}[\eta X - Y] &\geq P_0(\eta Q \alpha - 1)H - C_d m^{2d-2} \\ &\geq 3P_0 H - C_d m^{2d-2} \geq 5cH \geq cH. \end{aligned} \quad (\text{R23})$$

Some realization B' thus satisfies

$$\eta X - Y \geq cH. \quad (\text{R24})$$

There is no omitted class of tuples here: X and Y include all additive tuples in B' , with repeated entries counted. Hence

$$T_d(B') = X + Y \geq X \geq \frac{cH}{\eta} = \kappa_d \alpha^{2d} \eta^{2d-1} H, \quad (\text{R25})$$

and

$$\frac{T_d^\phi(B')}{T_d(B')} = \frac{X}{X+Y} \geq \frac{1}{1+\eta} \geq 1-\eta. \quad (\text{R26})$$

This proves both conclusions. $\square$

The weighted estimate is what makes the size condition depend on $(\alpha\eta)^{-2d}$. Counting all coefficient vectors separately would introduce a power of the bandwidth q . Instead, each unwanted relation is bounded by the same m^{2d-2} , and its contribution is multiplied by its Fourier weight; the weights have total mass one. The separate error for repeated entries has the same lower power of m . No independence assumption is made for repeated coordinates.

The exponent η^{2d-1} in the conclusion also deserves attention. The survival calculation supplies a quantity c proportional to $(\alpha\eta)^{2d}$, but (R24) yields $X \geq cH/\eta$. Keeping this final factor improves the retained count by η^{-1} .

II.9.4 Order eight and the comparison with Corollary 9.4

Corollary II.9.5 (Explicit order-eight restriction). *Let $0 < \beta, \gamma, \eta \leq 1$, let p be prime, and let $B \subset \mathbb{F}_p$ have size βp . Suppose that $\phi : B \rightarrow \mathbb{F}_p$ respects at least $\gamma |B|^3$ additive quadruples. If*

$$\boxed{p \geq 2^{102} \gamma^{-112} \eta^{-16} \beta^{-1}}, \quad (\text{R27})$$

then there is $B' \subset B$ containing at least

$$\boxed{2^{-95} \gamma^{112} \eta^{15} \beta^{15} p^{15}} \quad (\text{R28})$$

ordered additive 16-tuples, at least a proportion $1 - \eta$ of which are respected by ϕ .

Proof. Corollary II.8.2 permits $\alpha = \gamma^7$ in Theorem II.9.1 with $d = 8$. Its exact constants are

$$\kappa_8 = \frac{1}{2 \cdot 3200^8}, \quad \Lambda_8 = 242 \cdot 3200^8. \quad (\text{R29})$$

The integer inequalities

$$2 \cdot 3200^8 < 2^{95}, \quad 242 \cdot 3200^8 < 2^{102} \quad (\text{R30})$$

give the convenient powers of two in the statement. $\square$

The displayed count in source Corollary 9.4 is

$$\left(\frac{\gamma^7 \beta^6 \eta}{4} \right)^{2^{19}} \beta^{15} p^{15}. \quad (\text{R31})$$

The improvement has two distinct causes. Exact Parseval removes the β^6 inside the large power. The powered cosine replaces the repeated weak enrichment by one stronger selection, giving the powers $\gamma^{112} \eta^{15}$. Both changes preserve the same configuration normalization. The source corollary also inherits the sufficiently-large- p condition from Lemma 9.3; the present statement gives an explicit sufficient condition instead of suppressing it [1, 23].

This is a local improvement in the transition from a positive proportion of respected relations to an almost-homomorphism. Its relevance is not confined to the original exposition: later quantitative work also uses transitions of this kind, followed by additional steps that organize and extend the approximate homomorphism [6]. The present calculations do not by themselves supply an improved global bound for Szemerédi’s theorem. Such a claim would require propagating the parameters through all later restrictions and density increments.

[write] **6 October 2026.** Rechecked at the write: at $\gamma = \eta = \beta = 1/2$, $(\gamma^7 \beta^6 \eta / 4)^{2^{19}} = 2^{-16 \cdot 2^{19}} = 2^{-8\,388\,608}$, while $2^{-95} \gamma^{112} \eta^{15} = 2^{-222}$ relative to $|B|^{15}$, under $p \geq 2^{102} \cdot 2^{112} \cdot 2^{16} \cdot 2 = 2^{231}$ (Section V.16). The formal `lemma_9_3` and `corollary_9_4` (`Sections08_09.lean`, lines 77 and 91; exact companions in `Proofs09Restriction.lean`, lines 376 and 416) quantify over all moduli $N \geq N_0$ without primality and are proved by another route; Corollary II.9.5 is a prime-field interface and does not prove them, as Section V.19 says. Not formalized.

II.9.5 What the exponent says about this selection method

The use of a higher-degree kernel raises a natural optimization question. There is a simple bound that explains why the powers in Theorem II.9.1 are natural for the particular expectation argument just used.

Proposition II.9.6 (A limitation of single-kernel amplification). *Let K be a real trigonometric polynomial with $0 \leq K \leq 1$, mean $b_0 > 0$, and Fourier coefficients b_ℓ . For $d \geq 2$, put*

$$Q_K = \frac{\sum_\ell |b_\ell|^{2d}}{b_0^{2d}}.$$

Then

$$Q_K \leq \frac{1}{b_0}. \tag{R32}$$

Consequently, if $Q_K \geq 4/(\alpha\eta)$, the base survival probability b_0^{2d} is at most $(\alpha\eta/4)^{2d}$, and the scalar-multiple contribution $\sum_\ell |b_\ell|^{2d}$ to respected-tuple survival is at most $(\alpha\eta/4)^{2d-1}$.

Proof. Since $K \geq 0$, $|b_\ell| \leq b_0$. Parseval and $K^2 \leq K$ imply

$$\sum_\ell |b_\ell|^{2d} \leq b_0^{2d-2} \sum_\ell |b_\ell|^2 = b_0^{2d-2} \int_0^1 K(t)^2 dt \leq b_0^{2d-1}.$$

Divide by b_0^{2d} , and then use the assumed lower bound for Q_K . □

[write] **6 October 2026; sources 15–20.** Source 17’s Proposition II.62.2 determines the best retained count this scalar certificate can give, unconditionally, in terms of the respected fraction (its Lemma II.62.1 is (R32)).

For an input with exactly αH respected tuples, the contribution guaranteed by the scalar-multiple coefficient relations is bounded by a constant times $\alpha^{2d} \eta^{2d-1} H$, under the stipulated amplification requirement. Theorem II.9.1 attains these parameter powers up to constants depending on d . This is an observation about that part of the single-kernel expectation budget. Additional relations and repeated entries can increase the actual expectation. The calculation is not a lower bound for every possible restriction procedure, does not rule out unusually favorable realizations, and does not establish optimality of the theorem’s size threshold.

II.10 Arrangement moments retain the cross-section profile

Source 04, delivered Section 5.

The same principle applies to Lemmas 12.3 and 14.7, but their second moments carry multiplicities. The correct quantity to retain is the number of respected pairs of congruent cubes within a cross-section, rather than simply $|B|$.

II.10.1 The weighted representation

Let $B \subset \mathbb{F}_p^{k+1}$, with $k \geq 1$, let $\phi : B \rightarrow \mathbb{F}_p$, and write

$$B_r = \{y \in \mathbb{F}_p^k : (y, r) \in B\}.$$

For $h, y \in \mathbb{F}_p^k$, the axis-parallel cube with base y and sidelengths h has vertices

$$y + \epsilon \cdot h = (y_1 + \epsilon_1 h_1, \dots, y_k + \epsilon_k h_k), \quad \epsilon \in \{0, 1\}^k.$$

If all these vertices lie in B_r , define its value by

$$\Phi_h(y, r) = \sum_{\epsilon \in \{0, 1\}^k} (-1)^{\epsilon_1 + \dots + \epsilon_k} \phi(y + \epsilon \cdot h, r). \quad (\text{A1})$$

For each h , define the nonnegative integer weight

$$a_h(r, v) = \#\{y : y + \{0, 1\}^k \cdot h \subset B_r, \Phi_h(y, r) = v\}. \quad (\text{A2})$$

Let R_j count the respected j -arrangements: ordered lists of $2j$ cubes with common h , satisfying equal sums of the first and second halves of their cross-section indices, and equal sums of their cube values. Zero sidelengths and repeated vertices are included in this count, as in the moment identities of the source. Expanding the Fourier transform of a_h on $\mathbb{F}_p^2$ gives

$$R_j = \frac{1}{p^2} \sum_{h \in \mathbb{F}_p^k} \sum_{\chi \in \widehat{\mathbb{F}_p^2}} |\widehat{a}_h(\chi)|^{2j}. \quad (\text{A3})$$

In particular,

$$R_1 = \sum_{h, r, v} a_h(r, v)^2 \quad (\text{A4})$$

counts ordered congruent cube pairs in the same cross-section with equal values.

Theorem II.10.1 (Arrangement interpolation with profile). *Suppose that $R_2 > 0$. For every integer $j \geq 2$,*

$$R_j \geq \frac{R_2^{j-1}}{R_1^{j-2}}. \quad (\text{A5})$$

Define

$$\beta_r = \frac{|B_r|}{p^k}, \quad \beta = \frac{|B|}{p^{k+1}}, \quad \sigma = \frac{1}{p} \sum_r \beta_r^3. \quad (\text{A6})$$

Then

$$R_1 \leq \sigma p^{3k+1}, \quad \sigma \leq \beta. \quad (\text{A7})$$

Consequently, if $\theta > 0$ and $R_2 \geq \theta p^{5k+3}$, then

$$R_j \geq \theta^{j-1} \sigma^{-(j-2)} p^{(2j+1)k+2j-1}. \quad (\text{A8})$$

Proof. Apply the Hölder calculation of Theorem II.8.1 to the spectrum indexed by (h, χ) , using the common normalization p^{-2} in (A3). This gives (A5). The hypothesis $R_2 > 0$ ensures $R_1 > 0$, so the denominator is legitimate.

To bound R_1 , temporarily ignore the equality of cube values. A pair of cubes in B_r , with bases y, z and common sidelengths h , determines a triple

$$(y, z, y + h) \in B_r^3. \quad (\text{A9})$$

The last point is the opposite vertex of the first cube. This map is injective: the triple determines both bases and $h = (y + h) - y$. Hence there are at most $|B_r|^3$ such pairs, even before the value condition is imposed. Summing over r gives

$$R_1 \leq \sum_r |B_r|^3 = \sigma p^{3k+1}.$$

Since $0 \leq \beta_r \leq 1$, and $p^{-1} \sum_r \beta_r = \beta$, we have $\sigma \leq \beta$. Finally substitute these estimates into (A5). The ambient exponent is

$$(5k+3)(j-1) - (3k+1)(j-2) = (2j+1)k + 2j - 1,$$

which proves (A8). □

[write] 6 October 2026; sources 11–14. Source 14 classifies equality in this interpolation and in its profile bound (its Theorems II.44.1 and II.44.4), proves stability for the cross-section collision count, and improves the profile denominator from cross-section sizes (its Corollary II.44.3); (A13) and (A14) are special cases of its classification.

[write] 6 October 2026; sources 15–20. Source 15 proves (A5)–(A11) again, independently (its Theorem II.47.3, Lemma II.48.1, Corollary II.48.2), and couples the collision count with the product-property step: its Theorem II.49.6 gives $\beta^{22 \cdot 4^k} \gamma^{44k4^k + 56 \cdot 2^k}$ for eight-arrangements, against $\beta^{28 \cdot 4^k - 6} \gamma^{84k4^k}$ in (A12). The chain continues at Lemma 15.5 in source 18's Section II.66.

For $j = 8$, the replacement for Lemma 14.7 is therefore

$$R_8 \geq \theta^7 \sigma^{-6} p^{17k+15} \geq \theta^7 \beta^{-6} p^{17k+15}. \quad (\text{A10})$$

The source uses the same fourth-moment hypothesis but discards the factor σ when bounding the second moment. If every cross-section has density β , then $\sigma = \beta^3$, so the gain is β^{-18} . More generally, $\beta_r \leq K\beta$ for every r implies

$$\sigma \leq K^2 \beta^3. \quad (\text{A11})$$

This records exactly which information about the cross-sections can be used by the interpolation step.

Using Corollary 14.6 as stated in the repository's corrected transcription, the density power in Lemma 14.8 improves from $7 \cdot 4^{k+1}$ to $7 \cdot 4^{k+1} - 6$:

$$R_8 \geq \beta^{7 \cdot 4^{k+1} - 6} \gamma^{21k \cdot 4^{k+1}} p^{17k+15}. \quad (\text{A12})$$

Keeping the profile gives the stronger factor $\beta^{7 \cdot 4^{k+1}} \sigma^{-6}$. The analogous use in Lemma 12.4 changes $\beta^{112} \gamma^{336} p^{32}$ to $\beta^{106} \gamma^{336} p^{32}$. In that section the relevant fibers are vertical; its two coordinate roles are exchanged relative to the notation above. These conclusions use only the improved moment step, with the preceding hypotheses and estimates retained [1, 23].

II.10.2 Equality examples and the scope of sharpness

The arrangement interpolation and the profile bound extend to axis-parallel cubes over any finite abelian group G ; the proof uses only its characters and additive operations. This extension supplies a sharp example for the factor β^{-6} . Let $|G| = N$, let $C \leq G$, and take

$$B = G^k \times C, \quad \phi = 0, \quad \beta = \frac{|C|}{N}.$$

The nonempty cross-sections are full. Direct counting gives

$$R_1 = \beta N^{3k+1}, \quad R_2 = \beta^3 N^{5k+3}, \quad R_j = \beta^{2j-1} N^{(2j+1)k+2j-1}. \quad (\text{A13})$$

Thus both the interpolation inequality and $\sigma \leq \beta$ are equalities. This proves sharpness for the finite-abelian formulation; it does not assert the existence of proper subgroups of a cyclic group of prime order.

The uniform-profile case also has equality examples over a prime field. Let V be a coordinate subspace of $\mathbb{F}_p^k$, and take $B = V \times \mathbb{F}_p$, with $\phi = 0$. Every cross-section has density $\beta = |V|/p^k$. A cube is contained in V precisely when its base and sidelength vector belong to V . Therefore

$$R_j = |V|^{2j+1} p^{2j-1}, \quad \sigma = \beta^3. \quad (\text{A14})$$

Again the moment and profile estimates are equalities. In particular, the improvement by β^{-18} in the uniform-profile order-eight bound is attained at the available densities of these examples.

These equality constructions identify which parts of the argument cannot be improved using only the displayed moments and profile. They leave open the use of additional geometric or distributional information in applications.

II.11 Source 07's chapter: transport rigidity at Lemma 10.9 [write]

Added in the second write, 6 October 2026.

The next eight sections are source 07, a manuscript devoted to one step of Gowers's Section 10: the passage, in Lemma 10.9, from a local model of differences to a Freiman homomorphism. Its spine, an approximate homomorphism made exact, is shared with source 02's extension of Proposition 6.1, source 03's Sections II.5–II.6, source 04's restriction theorem and source 10's fibre extraction (Section II.21), but no other source treats Lemma 10.9, and source 07 shares with them only the notion of Freiman order and a reference on carries. It is therefore printed as a chapter of its own. Its notation table (its Appendix A) comes first, then its delivered Sections 2–8 in order: the fibered model and the directional error law (Section II.13), transport profiles (Section II.14), the coupled-path theorem (Section II.15), sharpness of $\delta < 1/(2k)$ (Section II.16), the application to Section 10 (Section II.17), the linear program for the starting law (Section II.18) and further extensions (Section II.19). Its introduction is Section F.12; its verification, questions, conclusion and reproduction commands are in Part V (Sections V.28–V.31).

Reading conventions. In this chapter δ is a failure probability (a directional defect), not a density; κ is a ratio of fibre sizes; $\mathcal{T}_s$ is a transport profile and $\mathcal{C}$ the cost of a word (not the cube count of Part III); $\theta, \eta, \alpha, \rho, \sigma$ and M are the parameters of Gowers's Lemmas 10.6–10.9 as recorded in `Section10.lean`. The table at the head of this Part lists the clashes with the rest of the report.

What is proved, in one paragraph. A violated relation between two words in the directions must consume total failure probability at least one along a pair of coupled paths with shared endpoints (Theorem II.15.1); with fibre-size ratio at most κ this forces a Freiman homomorphism of order k once $2[1 + (k-1)\kappa]\delta < 1$ (Corollary II.15.3), and for constant fibres the threshold $\delta < 1/(2k)$ cannot be increased (Theorem II.16.1). Under the hypotheses of Gowers's Lemma 10.9 this gives order five (Corollary II.17.3), and order nine when the construction's ambient fibre invariance, the fibre lower bound $R_g \geq \alpha^2 M/16$ and the parameter relations (II.17.10) are kept, none of which is a hypothesis of `lemma_10_9` (Corollary II.17.6). (Corrected after the independent check of 6 October

2026; the first wording read “and order nine when the ambient fibre invariance of the construction is kept”, which names only one of the three requirements.) Five and nine are guaranteed orders, not maximal ones, and the sharpness concerns δ , not η (the source’s own non-claims, Section F.12).

Relation to the formal statement (checked at HEAD 764740f07). The Lean proposition `lemma a_10_9` (`Section10.lean`, line 335; exact companion `lemma_10_9_holds`, `Proofs10Shift.lean`, line 717) concludes `FreimanHom 2 B' psi`. Corollary II.17.3 has exactly its hypotheses and concludes order five; Corollary II.17.6 needs in addition the fibre lower bound, the ambient invariance and the parameter relations (II.17.10), which in `Section10.lean` are hypotheses of `lemma_10_6` and of the Section 10 set-up but not of `lemma_10_9`. The precise comparison is in the notes after Corollaries II.17.3 and II.17.6. Neither statement is formalized.

II.12 Notation and assumptions at a glance

Source 07, delivered Appendix A.

Symbol	Meaning
G, H	Finite abelian index group and arbitrary abelian value group; Section II.19 also treats nonabelian groups.
X_g, R_g	Ambient fiber at index g and its cardinality; ambient fibers may be empty.
S, W, Ω	Nonempty usable support, union of its full fibers, and $ W $. Every R_g with $g \in S$ is positive.
$\mu(g)$	R_g/Ω on S , zero elsewhere. It is the index law of a uniform source point.
D, ϕ, ψ	Allowed directions, point-value map, and proposed directional difference map.
$F_d(g), \delta_d, b_d$	Conditional failure, its μ -mean, and source mass leaving S . Exits always count as failures.
$\mathcal{T}_s(t)$	Largest translated mass of a bounded error function of original mean at most t .
$c(s), A(s), M(s)$	Likelihood-ratio bound, positive mass discrepancy, and translated internal mass.
κ	Uniform ratio bound for fiber sizes on S .
L, r_0, τ	Ambient one-step fiber-size variation, internal lower bound, and $\tau = L/r_0$.
θ, η	Nested agreement error and overlap error in the Section 10 application.
$q(\theta), \Delta$	$2\theta - \theta^2$ and the convenient budget $q(\theta) + \eta$.
ν	An alternative initial-index law, supported on S ; the local error information remains measured using μ .

II.13 Finite fibered domains and the correct error law

Source 07, delivered Section 2.

Let G be a finite abelian group and H an abelian group, not necessarily finite. For each $g \in G$, let X_g be a finite set. Assume these sets are pairwise disjoint, and put $R_g = |X_g|$. Empty fibers are allowed. Fix a nonempty set

$$S \subseteq \{g \in G : R_g > 0\}, \quad W = \bigsqcup_{g \in S} X_g, \quad \Omega = |W|.$$

Let $\phi : \bigsqcup_g X_g \rightarrow H$, let $D \subseteq G$, and let $\psi : D \rightarrow H$ be the candidate difference map. The source-index law is

$$\mu(g) = \begin{cases} R_g/\Omega, & g \in S, \\ 0, & g \notin S. \end{cases} \quad (\text{II.13.1})$$

Thus choosing a point uniformly from W is the same as choosing its index with law μ , then choosing uniformly within that fiber.

Definition II.13.1 (Directional defect). For $d \in D$ and $g \in S$, set

$$F_d(g) = \begin{cases} 1, & g + d \notin S, \\ \frac{\#\{(x, y) \in X_g \times X_{g+d} : \phi(y) - \phi(x) \neq \psi(d)\}}{R_g R_{g+d}}, & g + d \in S. \end{cases} \quad (\text{II.13.2})$$

The total defect and boundary mass are

$$\delta_d = \sum_{g \in S} \mu(g) F_d(g), \quad b_d = \sum_{\substack{g \in S \\ g+d \notin S}} \mu(g). \quad (\text{II.13.3})$$

Operationally, start with a uniform point $x \in W$. A step in direction d is an automatic failure if its destination index lies outside S . Otherwise independently choose a uniform point y in the destination fiber and test the difference. This is a genuine probability experiment, and $0 \leq b_d \leq \delta_d \leq 1$.

An equivalent counting formula is

$$\delta_d = b_d + \frac{1}{\Omega} \sum_{\substack{g \in S \\ g+d \in S}} \frac{\#\{(x, y) \in X_g \times X_{g+d} : \phi(y) - \phi(x) \neq \psi(d)\}}{R_{g+d}}. \quad (\text{II.13.4})$$

The cancellation of R_g is worth noting. Uniform source points, uniform indices, and uniform ordered pairs across all fibers give different statistics when fiber sizes vary. All subsequent bounds use (II.13.3), not an unnormalized pair-counting surrogate.

Declaring a departure from W to be a failure is conservative: some outside fibers may exist and still carry correct differences. The advantage is that no later step depends on an unproved assertion about those fibers. In particular, we never divide by the size of an empty fiber.

Definition II.13.2 (Freiman order). For $k \geq 1$, a map $\psi : D \rightarrow H$ is a Freiman homomorphism of order k if, for all $d_1, \dots, d_k, e_1, \dots, e_k \in D$,

$$\sum_{i=1}^k d_i = \sum_{i=1}^k e_i \implies \sum_{i=1}^k \psi(d_i) = \sum_{i=1}^k \psi(e_i).$$

This definition makes sense for an empty direction set, when the assertion is vacuous. If $D \neq \emptyset$, preservation of order k implies preservation of every smaller positive order: pad both sides with repetitions of one fixed element of D , then cancel their images.

II.14 Transport profiles instead of square-root losses

Source 07, delivered Section 3.

II.14.1 An exact finite optimization

For $s \in G$, the shifted law is $\mu_s(g) = \mu(g - s)$. We only need its restriction to S ; mass outside S corresponds to discarded locations.

Definition II.14.1 (Transport profile). For $0 \leq t \leq 1$, define

$$\mathcal{T}_s(t) = \max \left\{ \sum_{g \in S} \mu(g-s)a_g : 0 \leq a_g \leq 1, \quad \sum_{g \in S} \mu(g)a_g \leq t \right\}. \quad (\text{II.14.1})$$

Also define

$$c(s) = \max_{g \in S} \frac{\mu(g-s)}{\mu(g)}, \quad A(s) = \sum_{g \in S} (\mu(g-s) - \mu(g))_+, \quad M(s) = \sum_{g \in S} \mu(g-s). \quad (\text{II.14.2})$$

Here $x_+ = \max\{x, 0\}$. Every denominator in (II.14.2) is positive. The feasible set in (II.14.1) is nonempty and compact, so the maximum exists. The quantity $\mathcal{T}_s(t)$ is the largest translated mass that an error function of original mean at most t can acquire.

Proposition II.14.2 (Basic profile bounds). *The profile is nondecreasing, concave, and piecewise linear. Moreover,*

$$\mathcal{T}_0(t) = t, \quad \mathcal{T}_s(t) \leq \min\{c(s)t, t + A(s), M(s)\}. \quad (\text{II.14.3})$$

An exact evaluation is obtained by sorting the ratios $\lambda_g = \mu(g-s)/\mu(g)$ in decreasing order and filling original μ -mass up to t , using at most one fractional final coordinate.

Proof. Increasing the budget enlarges the feasible set. If a and a' are feasible for budgets t and t' , their convex combination is feasible for the corresponding convex combination of the budgets; the objective is linear. Taking maximizing vectors proves concavity.

For any feasible a ,

$$\sum_g \mu(g-s)a_g \leq c(s) \sum_g \mu(g)a_g \leq c(s)t.$$

Subtracting $\sum_g \mu(g)a_g$ and retaining only positive coefficient differences gives the bound $t + A(s)$. The condition $a_g \leq 1$ gives $M(s)$. For $s = 0$, the objective equals the constraint sum, and the constant vector $a_g = t$ attains t .

For the exact formula, use variables $x_g = \mu(g)a_g$. The problem becomes

$$\max \sum_g \lambda_g x_g, \quad 0 \leq x_g \leq \mu(g), \quad \sum_g x_g \leq t.$$

If $\lambda_g > \lambda_h$, $x_g < \mu(g)$, and $x_h > 0$, moving a small positive mass from h to g increases the objective. Repeating this exchange leaves all higher ratios filled before a lower one is used. Ties may be resolved arbitrarily. This proves the stated algorithm and its piecewise linearity, including zero-ratio coordinates. $\square$

Example II.14.3 (The profile can beat the coarse envelopes). On $G = \mathbb{Z}/2\mathbb{Z}$, take $S = G$ and $\mu = (1/3, 2/3)$. For the nonzero shift,

$$\mathcal{T}_1(t) = \begin{cases} 2t, & 0 \leq t \leq 1/3, \\ \frac{1}{2} + \frac{1}{2}t, & 1/3 \leq t \leq 1. \end{cases}$$

Here $c(1) = 2$ and $A(1) = 1/3$. At $t = 1/2$, the exact value is $3/4$, whereas $\min\{2t, t + 1/3, 1\} = 5/6$.

II.14.2 Comparison and additive invariance of fiber sizes

Lemma II.14.4 (Uniform and prefixwise transport). *Suppose $R_u/R_v \leq \kappa$ for all $u, v \in S$, where $\kappa \geq 1$. Then $c(s) \leq \kappa$ for every $s \in G$.*

Suppose in addition that, for some $r_0 > 0$, $L \geq 0$,

$$R_g \geq r_0 \quad (g \in S), \quad |R_{g+d} - R_g| \leq L \quad (g \in G, d \in D). \quad (\text{II.14.4})$$

Put $\tau = L/r_0$. If $s = d_1 + \dots + d_j$ with $d_i \in D$, then

$$c(s) \leq 1 + j\tau. \quad (\text{II.14.5})$$

When both hypotheses hold, the minimum of κ and $1 + j\tau$ is available.

Proof. If $g - s \notin S$, the ratio defining $c(s)$ is zero. Otherwise it equals R_{g-s}/R_g , which gives the first assertion. For the second, telescope (II.14.4) along the sequence from $g - s$ to g . All intermediate indices are allowed because invariance is assumed on all of G , not just on S . This gives $R_{g-s} \leq R_g + jL$, and division by $R_g \geq r_0$ proves the claim. $\square$

The uniform estimate does not become κ^j after j steps. We compare the translated law with the original law directly, rather than renormalizing the surviving sample at each step. This is one reason the path argument retains useful higher-order information.

II.14.3 A parameterized replacement for a nested shift estimate

Proposition II.14.5 (Average transport of a point property). *Let $Q \subseteq W$, with $|W \setminus Q|/\Omega \leq \varepsilon$. For an internal destination fiber put $a_g = |X_g \setminus Q|/R_g$. Start at an index with law μ and shift by d . The mean destination failure restricted to destinations in S is at most $\mathcal{T}_d(\varepsilon)$.*

Consequently, for every $u > 0$, a proportion at least

$$1 - b_d - \frac{\mathcal{T}_d(\varepsilon)}{u} \quad (\text{II.14.6})$$

of source points have a destination in S on whose fiber the fraction outside Q is at most u .

Proof. The original mean of a is $|W \setminus Q|/\Omega \leq \varepsilon$. The destination mean over internal shifts is

$$\sum_{\substack{g \in S \\ g+d \in S}} \mu(g) a_{g+d} = \sum_{h \in S} \mu(h-d) a_h \leq \mathcal{T}_d(\varepsilon).$$

Markov's inequality bounds the source mass of internal destinations with $a_{g+d} > u$ by this mean divided by u . Add the boundary mass b_d . $\square$

A square root appears if one insists on making the outer exceptional fraction and the inner failure threshold comparable. It is not intrinsic to the averaged statement. The next section uses the average itself, so no such balancing is necessary.

II.15 The coupled-path rigidity theorem

Source 07, delivered Section 4.

For a nonempty word $\mathbf{d} = (d_1, \dots, d_m)$ in D , define its prefixes $s_0 = 0$, $s_j = d_1 + \dots + d_j$, and its cost

$$\mathcal{C}(\mathbf{d}) = \sum_{j=1}^m \mathcal{T}_{s_{j-1}}(\delta_{d_j}). \quad (\text{II.15.1})$$

Theorem II.15.1 (Transport-sensitive polygon rigidity). *Let $\mathbf{d} = (d_1, \dots, d_m)$ and $\mathbf{e} = (e_1, \dots, e_n)$ be nonempty words in D , with $\sum_j d_j = \sum_j e_j$. If*

$$\mathcal{C}(\mathbf{d}) + \mathcal{C}(\mathbf{e}) < 1, \quad (\text{II.15.2})$$

then

$$\sum_{j=1}^m \psi(d_j) = \sum_{j=1}^n \psi(e_j). \quad (\text{II.15.3})$$

Equivalently, every relation violated by ψ has total cost at least one. Each profile in the hypothesis may be replaced by any valid upper bound from Proposition II.14.2.

Proof. We specify the entire finite experiment before using a union bound. Choose $r \in S$ with law μ . The first path has formal indices $r + s_0, \dots, r + s_m$, and the second has formal indices $r + t_0, \dots, r + t_n$, where $t_j = e_1 + \dots + e_j$. Their initial indices are the same, and their terminal indices are the same.

At the common initial position choose one uniform point $x \in X_r$. If the terminal index belongs to S , choose one uniform terminal point y in that fiber, independently of x . Use this same y for both paths. At all internal formal positions whose indices lie in S , choose uniform fiber points independently of each other, of x , and of y . Different formal positions remain independent even when their indices coincide. In particular, when the total shift is zero, the initial and terminal points are still two independent draws from the same fiber. Positions outside S receive dummy values and will be treated as failures.

For edge j on the first path, let E_j be the event that all the preceding indices $r + s_0, \dots, r + s_{j-1}$ lie in S , and either the next index leaves S , or its edge equation fails. “Preceding indices” means only index survival; it does *not* require preceding equations to hold. Thus this restriction depends only on r , not on the sampled fiber values.

Write $g = r + s_{j-1}$. For a surviving source index g , the conditional failure probability is exactly $F_{d_j}(g)$: its endpoint samples have the required independent uniform marginal laws. The mass of the possible starting index $r = g - s_{j-1}$ is $\mu(g - s_{j-1})$. Discarding the extra requirement that all earlier indices survive therefore gives


$$\begin{aligned} \mathbb{P}(E_j) &\leq \sum_{g \in S} \mu(g - s_{j-1}) F_{d_j}(g) \\ &\leq \mathcal{T}_{s_{j-1}}(\delta_{d_j}). \end{aligned} \quad (\text{II.15.4})$$

The last step uses $0 \leq F_{d_j} \leq 1$ and its original mean δ_{d_j} . The same construction for the second path gives events E'_j with the corresponding bounds.

If none of these events occurs, no path can leave S : its first exit would be one of the events. Every edge equation then holds. Telescoping along both paths yields

$$\sum_j \psi(d_j) = \phi(y) - \phi(x) = \sum_j \psi(e_j).$$

The union bound and (II.15.2) show that the probability of avoiding all events is positive, proving the asserted equality. Alternatively, if the equality fails, the union of the events is the whole sample space, so the sum of their upper bounds is at least one. $\square$



One shared start and one shared terminal; the two intermediate draws are independent, including when their fiber indices coincide.

For this four-edge diagram the criterion reads explicitly

$$\delta_a + \mathcal{T}_a(\delta_b) + \delta_c + \mathcal{T}_c(\delta_d) < 1. \quad (\text{II.15.5})$$

There is no independence assumption between edge failures. The union bound needs only their marginals. Conversely, conditioning each successive sample on all earlier equations being correct would change those marginals and would not justify (II.15.4).

Remark II.15.2 (Path-specific survival can be retained). For an edge with prefix s_{j-1} , the exact source-index weight in the proof is

$$\mu(g - s_{j-1}) \prod_{i=0}^{j-1} \mathbf{1}_S(g - s_{j-1} + s_i), \quad g \in S.$$

Replacing the translated weight in that edge's profile by this smaller weight gives another valid criterion. The simpler translation-only profile is useful because it depends on the prefix, not the full path history.

II.15.1 Higher-order consequences

Corollary II.15.3 (Explicit Freiman criteria). *Suppose $\delta_d \leq \delta$ for all $d \in D$, with $0 \leq \delta \leq 1$, and $k \geq 2$.*

If the fiber ratio on S is at most κ , then

$$2[1 + (k-1)\kappa]\delta < 1 \implies \psi \in \text{Fr}_k(D, H). \quad (\text{II.15.6})$$

Under (II.14.4),

$$[2k + \tau k(k-1)]\delta < 1 \implies \psi \in \text{Fr}_k(D, H). \quad (\text{II.15.7})$$

If both forms of control are available, it is enough that

$$2\delta \left[1 + \sum_{j=1}^{k-1} \min\{\kappa, 1 + j\tau\} \right] < 1. \quad (\text{II.15.8})$$

Proof. For a word of length k , the first cost is at most δ , because $\mathcal{T}_0(t) = t$. Every other cost is at most $\kappa\delta$, proving a one-path bound of $[1 + (k-1)\kappa]\delta$. Alternatively, Lemma II.14.4 bounds the $(j+1)$ -st cost by $(1 + j\tau)\delta$. Summing over $j = 0, \dots, k-1$ gives $[k + \tau k(k-1)/2]\delta$. Double the applicable one-path bound and use Theorem II.15.1 for an arbitrary order- k relation. Taking the smaller bound at each prefix proves (II.15.8). $\square$

[write] **6 October 2026; sources 34–39.** *Sharp for every κ* (source 34, Appendix A). In the finite-fibre model of this section the threshold $1/(2[1 + (k-1)\kappa]) = 1/(2 + (2k-2)\kappa)$ is optimal for every $k \geq 2$ and every $\kappa \geq 1$, already with zero boundary loss (Theorem II.163.1); this answers source 07’s Question 1 (item II.9). The formal lemma_10_9 concerns the prime-order interface and is unaffected.

In the important factor-two case, (II.15.6) is

$$(4k-2)\delta < 1. \quad (\text{II.15.9})$$

When the fibers have constant size on S , the sufficient condition is $2k\delta < 1$, even when S is a proper subset of G . Any boundary cost is already included in δ .

II.16 The balanced threshold is sharp

Source 07, delivered Section 5.

Theorem II.16.1 (Sharp uniform directional threshold). *For every $k \geq 2$, the universal strict threshold $\delta < 1/(2k)$ in the constant-fiber setting cannot be increased. More precisely, there are examples with full support, singleton fibers, identical finite cyclic source and target groups, and $\delta_d = 1/(2k)$ for every selected direction, whose candidate map is not a Freiman homomorphism of order k .*

Proof. Choose any integer $m \geq 2$ and put $N = 2km$. Take $G = H = \mathbb{Z}/N\mathbb{Z}$, $S = G$, and $X_r = \{r\}$. Using the representatives $0 \leq r < N$, define

$$\phi(r) = \left\lfloor \frac{r}{m} \right\rfloor \pmod{N}, \quad D = \{m, -m\}, \quad \psi(m) = 1, \quad \psi(-m) = -1. \quad (\text{II.16.1})$$

The two directions are distinct because $k \geq 2$.

For $0 \leq r < N - m$, shifting by m increases $\lfloor r/m \rfloor$ by one. On the final block of m residues, wraparound changes the value by $-(2k-1)$. This is not congruent to one modulo N , since their difference is $2k$ and $N = 2km > 2k$. Thus exactly m of the N source points fail, and $\delta_m = m/N = 1/(2k)$. Shifting by $-m$ has the analogous failure on the first block, so $\delta_{-m} = 1/(2k)$.

In the source group the k -term sums satisfy

$$\underbrace{m + \dots + m}_{k \text{ terms}} = \underbrace{(-m) + \dots + (-m)}_{k \text{ terms}},$$

because their difference is $2km = N$. Their proposed image sums are k and $-k$, whose difference $2k$ is nonzero modulo N . Therefore the order- k relation is violated. Its two-path budget is exactly $2k\delta = 1$, showing why strict inequality is necessary. $\square$

[write] **6 October 2026; sources 34–39.** Source 34’s Theorem II.163.1 extends the sharpness to every fibre ratio κ (note after Corollary II.15.3).

Example II.16.2 (An eight-point obstruction). For $k = m = 2$, the values of ϕ on $\mathbb{Z}/8\mathbb{Z}$ are

$$(0, 0, 1, 1, 2, 2, 3, 3).$$

With directions 2, 6 and labels 1, -1 , both defect fractions are $1/4$. Although $2 + 2 = 6 + 6$ modulo eight, the two image sums are 2 and 6. The domain has no missing fibers and no nonuniformity; the obstruction is a genuine accumulated wraparound error.

The example is naturally described as a carry or winding obstruction; carry phenomena and approximate homomorphisms also appear in [28]. What is established here is the exact extremal statement in Theorem II.16.1, with its specific directional sampling model.

The restriction to arbitrary selected direction sets matters. For example, if $D = G$, the constant-fiber criterion applied to the words (a, b) and $(a + b)$ shows that a uniform bound $3\delta < 1$ already forces $\psi(a + b) = \psi(a) + \psi(b)$. In particular, the sharp $1/4$ threshold for arbitrary-direction order-two relations is *not* a sharp threshold for full-group homomorphism recovery.

II.17 Application to Gowers’s Section 10

Source 07, delivered Section 6.

II.17.1 The local input and its averaged consequence

The following formulation isolates the assumptions used from the repository interface [42]. Let $B \subseteq G$ be symmetric and $D = B' \subseteq B$. Suppose that W is a nonempty union of entire fibers, that their sizes on W vary by at most a factor of two, and that

$$|W \cap (W + d)| \geq (1 - \eta)|W| \quad (d \in B). \quad (\text{II.17.1})$$

Here $W + d$ denotes the union of the ambient fibers with indices in $S + d$; it is not a chosen bijection between individual fiber points. For each $d \in D$, suppose at least a fraction $1 - \theta$ of points $x \in W$ satisfy

$$\#\{y \in X_{r(x)+d} : \phi(y) - \phi(x) = \psi(d)\} \geq (1 - \theta)|X_{r(x)+d}|, \quad (\text{II.17.2})$$

where $r(x)$ denotes the index of x . For now assume $0 \leq \theta \leq 1$ and $\eta \geq 0$.

Because $-d \in B$, (II.17.1) implies

$$b_d = 1 - \frac{|W \cap (W - d)|}{|W|} \leq \eta. \quad (\text{II.17.3})$$

This is the precise point at which symmetry is used. The main path theorem itself does not require a symmetric direction set.

Lemma II.17.1 (Nested agreement to a single defect budget). *Under (II.17.2), put $q(\theta) = 2\theta - \theta^2$. Then*

$$\delta_d \leq 1 - (1 - \theta)(1 - \theta - b_d)_+ = \min\{1, q(\theta) + (1 - \theta)b_d\}. \quad (\text{II.17.4})$$

In particular, under (II.17.3),

$$\delta_d \leq q(\theta) + \eta. \quad (\text{II.17.5})$$

Proof. The good-source set in (II.17.2) has source-point mass at least $1 - \theta$. The set of sources with destination in S has mass $1 - b_d$. Their intersection has mass at least $(1 - \theta - b_d)_+$. On this intersection the destination fiber is nonempty, and the conditional success probability is at least $1 - \theta$. Therefore the overall success probability $1 - \delta_d$ is at least their product. Expanding gives the minimum formula, and $b_d \leq \eta$ gives (II.17.5). $\square$

This proof includes the case where (II.17.2) is vacuous on an empty ambient target fiber. Such a source lies outside the internal intersection and is charged to b_d . There is no hidden nonemptiness assumption.

For readable parameter comparisons we use the slightly larger budget

$$\Delta = q(\theta) + \eta. \quad (\text{II.17.6})$$

Every assertion below remains valid with the smaller budget from (II.17.4). Whenever a displayed criterion has the form $C\Delta < 1$ with $C \geq 1$, it automatically ensures $\Delta < 1$, as required when it is substituted for δ in the preceding results.

Theorem II.17.2 (A replacement exactification interface). *Under the local assumptions (II.17.1)–(II.17.2) and the factor-two fiber bound, for every $k \geq 2$,*

$$(4k - 2)(2\theta - \theta^2 + \eta) < 1 \implies \psi \in \text{Fr}_k(D, H). \quad (\text{II.17.7})$$

If ambient invariance (II.14.4) also holds, the alternative sufficient criterion is

$$[2k + \tau k(k - 1)](2\theta - \theta^2 + \eta) < 1. \quad (\text{II.17.8})$$

Proof. Use Lemma II.17.1, followed by Corollary II.15.3. No approximate-additivity or energy hypothesis beyond those already used to construct the local model is invoked here. $\square$

The density lower bound on B' in the repository's local-model definition is not needed by this exactification theorem. Its relevance lies in obtaining a large domain for later arguments, not in verifying a relation once its participating directions belong to B' .

II.17.2 Order five under the recorded Lemma 10.9 hypotheses

Corollary II.17.3 (Order five with no stronger smallness condition). *Under the local factor-two assumptions above, assume $\theta = 10\eta^{1/5}$ and $6\sqrt{\theta} < 1$. Then $\psi \in \text{Fr}_5(D, H)$.*

Proof. The assumptions give

$$\theta < \frac{1}{36}, \quad \eta < 360^{-5}.$$

Since q is increasing on $[0, 1]$,

$$q(\theta) < q(1/36) = \frac{71}{1296}.$$

The factor in (II.17.7) for $k = 5$ is eighteen. Hence

$$18\Delta < \frac{71}{72} + \frac{18}{360^5} < 1,$$

because $1296 < 360^5$. Apply Theorem II.17.2. $\square$

For the literal Lean statement there is a minor sign convention to address. If B' is empty, the order-five conclusion is vacuous. Otherwise choose $d \in B' \subseteq B$; nonemptiness of W and (II.17.1), together with $|W \cap (W + d)| \leq |W|$, force $\eta \geq 0$. Thus introducing a nonnegative real fifth root does not add an assumption in the nonvacuous case. The existing proof also derives this sign fact from overlap [43].

[write] 6 October 2026; sources 06–10. *Exactly the Lean hypotheses (checked at HEAD 764740f07).* lemma_10_9 (Section10.lean, line 335) assumes W nonempty, IsSection10ShiftRegular D W B eta (fibre saturation of W , symmetry of B , fibre sizes on W within a factor two, and $(1 - \eta)|W| \leq |W \cap D.\text{shift } W \ d|$ for $d \in B$), IsSection10LocalDifferenceModel D ϕ W B B' ψ θ with $\theta = 10\eta^{1/5}$ ($B' \subseteq B$, $(1 - \theta)|B| \leq |B'|$, and for $d \in B'$ the nested statement that a proportion $1 - \theta$ of $w \in W$ have a proportion $1 - \theta$ of z in the fibre over $\text{index}(w) + d$ with $\phi(z) - \phi(w) = \psi(d)$), and $6\sqrt{\theta} < 1$; it concludes FreimanHom 2 B' ψ , the Mathlib notion of the definition in Section II.13 with $G = H = \mathbb{Z}/N\mathbb{Z}$. These are the assumptions of Corollary II.17.3: the overlap is (II.17.1), the nested condition is (II.17.2), and the density condition on B' is not used (paragraph after Theorem II.17.2). So Corollary II.17.3 is the statement of lemma_10_9 with FreimanHom 5 in place of FreimanHom 2. It is not formalized; the companion lemma_10_9_holds (Proofs10Shift.lean, line 717) proves order two, and both files are unchanged since the pin ae0e89574.

Proposition II.17.4 (A larger allowance at the same factor-two interface). *Retain the local factor-two assumptions and $\theta = 10\eta^{1/5}$. The condition*

$$0 \leq \eta \leq 120^{-5} \tag{II.17.9}$$

is sufficient for $\psi \in \text{Fr}_2(D, H)$.

Proof. Now $\theta \leq 1/12$, so $q(\theta) \leq 23/144$. Therefore

$$6\Delta \leq \frac{23}{24} + \frac{6}{120^5} < 1.$$

The case $k = 2$ of (II.17.7) proves the assertion. □

The ratio of this explicit allowance to 360^{-5} is $(360/120)^5 = 243$. This compares sufficient conditions for the local step only; it does not assert a factor-243 improvement in any final progression-length or density estimate.

II.17.3 Retaining the full ambient invariance

The construction in Section 10 supplies stronger fiber control than the minimal shift-regular interface records. For its application here, the needed inequalities are the following [42]:

$$\begin{aligned} R_g &\geq \alpha^2 M / 16 \quad (g \in S), & |R_{g+d} - R_g| &\leq \sigma M \quad (g \in G, d \in B), \\ 0 < \alpha &\leq 1, \quad M > 0, \quad 0 \leq \rho &\leq \alpha^2 / 32, \quad 0 \leq \sigma &\leq \eta \rho \alpha^2 / 16. \end{aligned} \tag{II.17.10}$$

The relevant transport parameter is consequently

$$\tau = \frac{16\sigma}{\alpha^2} \leq \eta \rho \leq \eta. \tag{II.17.11}$$

Only these displayed consequences, not a new construction of W , are used in this subsection.

Theorem II.17.5 (An explicit higher-order range). *Assume the local model, the factor-two regularity, and (II.17.10). Set $\theta = 10\eta^{1/5}$. For every integer $k \geq 2$,*

$$0 \leq \eta \leq (40k)^{-5} \implies \psi \in \text{Fr}_k(D, H). \quad (\text{II.17.12})$$

More precisely, in the range $\eta \geq 0$ and $0 \leq \theta \leq 1$, it suffices that

$$[2k + \eta k(k-1)](2\theta - \theta^2 + \eta) < 1. \quad (\text{II.17.13})$$

Proof. The precise criterion follows from (II.17.11) and (II.17.8). To verify the convenient allowance (II.17.12), observe that

$$\theta \leq \frac{1}{4k}, \quad \Delta \leq \frac{1}{2k} - \frac{1}{16k^2} + \eta.$$

Using $k(k-1) \leq k^2$ and expanding gives

$$\begin{aligned} [2k + \eta k(k-1)]\Delta &\leq (2k + \eta k^2) \left(\frac{1}{2k} - \frac{1}{16k^2} + \eta \right) \\ &= 1 - \frac{1}{8k} + \frac{5}{2}k\eta - \frac{\eta}{16} + k^2\eta^2 \\ &\leq 1 - \frac{1}{8k} + \frac{5}{2}k\eta + k^2\eta^2. \end{aligned}$$

The hypothesis implies $\eta \leq 1/(64k^2)$, since $(40k)^5 \geq 64k^2$. Hence the last two positive terms are at most

$$\frac{5}{128k} + \frac{1}{4096k^2} \leq \frac{161}{4096k} < \frac{1}{8k}.$$

The strict inequality needed for (II.17.13) follows, including at the endpoint $\eta = (40k)^{-5}$. □

Corollary II.17.6 (Order nine in the original parameter range). *Under the full assumptions of Theorem II.17.5, the original condition $6\sqrt{10\eta^{1/5}} < 1$ implies $\psi \in \text{Fr}_9(D, H)$.*

Proof. It gives $\eta < 360^{-5} = (40 \cdot 9)^{-5}$, so use Theorem II.17.5 with $k = 9$. □

The distinction between Corollaries II.17.3 and II.17.6 is important. Order five follows from the assumptions exposed in the recorded local shift interface. Order nine additionally retains ambient invariance and a fiber lower bound from the construction. Omitting those extra hypotheses from an order-nine theorem would not be justified by this proof.

[write] **6 October 2026; sources 06–10.** *Order nine is not a corollary of the Lean statement.* Theorem II.17.5 and Corollary II.17.6 use, besides the hypotheses of `lemma_10_9` listed in the previous note, the package (II.17.10): fibre sizes at least $\alpha^2 M/16$ on S ; $|R_{g+d} - R_g| \leq \sigma M$ for all $g \in G$ and $d \in B$; $0 < \alpha \leq 1$, $M > 0$, $0 \leq \rho \leq \alpha^2/32$ and $0 \leq \sigma \leq \eta\rho\alpha^2/16$. None of these is a hypothesis of `lemma_10_9`. In `Section10.lean` they occur in other places: the fibre lower bound in `IsSection10RegularComponent`, the ambient invariance as `DomainInvariant D B (sigma * M)` and the bounds on α , M in `Section10Setup`, and $\rho \leq \alpha^2/32$, $\sigma \leq \eta\rho\alpha^2/16$ as hypotheses of `lemma_10_6` (there with $\sigma > 0$ and no explicit sign condition on ρ). A formal order-nine theorem therefore needs a new statement carrying these assumptions, for example one under the hypotheses of `lemma_10_6`, not a corollary of `lemma_10_9`. The order-five and order-nine statements are kept separate here, as the source keeps them.

II.17.4 A larger order-two allowance with ambient invariance

Proposition II.17.7 (The allowance 75^{-5}). *Under the full local and ambient assumptions of Theorem II.17.5, the condition*

$$0 \leq \eta \leq 75^{-5} \quad (\text{II.17.14})$$

is sufficient for $\psi \in \text{Fr}_2(D, H)$.

Proof. We have $\theta \leq 2/15$, so $q(\theta) \leq 56/225$. It follows that

$$\begin{aligned} (4 + 2\eta)\Delta &\leq (4 + 2\eta) \left(\frac{56}{225} + \eta \right) \\ &= \frac{224}{225} + \left(4 + \frac{112}{225} \right) \eta + 2\eta^2. \end{aligned}$$

Because $\eta \leq 75^{-5} < 1/4$, the last two terms are at most 5η . Indeed, $4 + 112/225 + 1/2 < 5$. Finally $5\eta \leq 5/75^5 < 1/225$, so the total is strictly below one. This is (II.17.13) for $k = 2$. $\square$

The two explicit endpoints and their ratio are

$$\begin{aligned} 360^{-5} &= \frac{1}{6046617600000} \approx 1.65382 \times 10^{-13}, \\ 75^{-5} &= \frac{1}{2373046875} \approx 4.21399 \times 10^{-10}, \\ \frac{75^{-5}}{360^{-5}} &= \left(\frac{24}{5} \right)^5 = 2548.03968. \end{aligned} \quad (\text{II.17.15})$$

These are transparent, rounded sufficient thresholds, not optimized roots. The exact inequalities (II.17.7) and (II.17.13), or the profiles themselves, should be used when constant-level optimization matters.

What does not follow. The earlier extraction of a large regular component, the size of B' , subsequent Bohr-set shrinkage, and the final density iteration all retain their own requirements. Enlarging the allowance at this one compatibility step does not remove them. In particular, neither the number 243 nor 2548.03968 is a claimed improvement factor in a global Szemerédi bound.

II.18 Choosing the initial law by a rational linear program

Source 07, delivered Section 7.

The original source law μ is useful because it is the law for which local errors are known. It need not be the best law from which to start a particular pair of paths. The distinction can be handled exactly.

Let ν be any probability law supported on S . Keep the same local functions F_d and their μ -means δ_d , but define

$$\mathcal{T}_s^\nu(t) = \max \left\{ \sum_{g \in S} \nu(g - s) a_g : 0 \leq a_g \leq 1, \quad \sum_{g \in S} \mu(g) a_g \leq t \right\}. \quad (\text{II.18.1})$$

Proposition II.18.1 (Alternative-root certificate). *Theorem II.15.1 remains valid after replacing every $\mathcal{T}_s$ by $\mathcal{T}_s^\nu$, provided that the same law ν is used for both paths. The choice of ν may depend on the relation being verified.*

Proof. Start the common initial index with law ν , still choosing uniformly within its fiber. The source-index weight for an edge of prefix s is bounded by $\nu(g - s)$, while F_d still has known μ -mean δ_d . Thus (II.15.4) becomes the corresponding bound with $\mathcal{T}_s^\nu$. All other steps are unchanged. $\square$

Lemma II.18.2 (An exact dual expression). *For every supported law ν , shift s , and $0 \leq t \leq 1$,*

$$\mathcal{T}_s^\nu(t) = \min_{\lambda \geq 0} \left[\lambda t + \sum_{g \in S} (\nu(g - s) - \lambda \mu(g))_+ \right]. \quad (\text{II.18.2})$$

Proof. For any feasible vector a and any $\lambda \geq 0$,

$$\begin{aligned} \sum_g \nu(g - s) a_g &= \lambda \sum_g \mu(g) a_g + \sum_g [\nu(g - s) - \lambda \mu(g)] a_g \\ &\leq \lambda t + \sum_g (\nu(g - s) - \lambda \mu(g))_+. \end{aligned}$$

This proves one inequality. For the reverse, order the ratios $\nu(g - s)/\mu(g)$ and use the filling construction in Proposition II.14.2. Choose λ to be the ratio of the possibly fractional final coordinate. Every positive coefficient in the second sum is then fully filled, every negative one is unfilled, and the remaining coordinate has zero coefficient. This makes the displayed bound an equality. At $t = 0$, choose λ at least the largest ratio; at $t = 1$, choose $\lambda = 0$. Budget boundaries and tied ratios give the same equality by the identical argument. $\square$

Theorem II.18.3 (Optimizing the root is a finite linear program). *Fix two equal-total words. Index all their edges by a finite set I , and let edge i have prefix s_i and direction d_i . The smallest alternative-root budget equals the optimum of*

$$\begin{aligned} &\text{minimize} \quad \sum_{i \in I} \left[\lambda_i \delta_{d_i} + \sum_{g \in S} z_{i,g} \right], \\ &\text{subject to} \quad \nu(g) \geq 0 \ (g \in S), \quad \sum_{g \in S} \nu(g) = 1, \quad \nu(g) = 0 \ (g \notin S), \\ &\quad \lambda_i \geq 0, \quad z_{i,g} \geq 0, \quad z_{i,g} \geq \nu(g - s_i) - \lambda_i \mu(g) \quad (i \in I, g \in S). \end{aligned} \quad (\text{II.18.3})$$

An optimum below one certifies the image relation. When the input laws and budgets are rational, an optimal rational certificate exists.

Proof. For fixed ν , independently minimizing over each λ_i and $z_{i,g}$ gives exactly the sum of the dual expressions (II.18.2). Minimizing also over ν therefore gives the desired smallest budget. Proposition II.18.1 proves the certificate claim.

For attainment and rationality, let $\mu_{\min} = \min_{g \in S} \mu(g) > 0$. One may restrict $0 \leq \lambda_i \leq 1/\mu_{\min}$ without worsening the objective: at the upper endpoint every coefficient $\nu(g - s_i) - \lambda_i \mu(g)$ is already nonpositive. Likewise one may restrict $0 \leq z_{i,g} \leq 1$, since the minimal feasible positive part never exceeds one. The resulting feasible set is a nonempty compact polytope. With rational data it is a rational polytope, and a linear objective attains its minimum at a rational vertex. $\square$

Taking $\nu = \mu$ recovers the earlier theorem, so this optimization never weakens that certificate. It is not claimed that every compatible relation admits a budget below one. The profiles optimize over arbitrary bounded error functions with the specified means; the errors of an actual map ϕ satisfy additional simultaneous constraints that this relaxation need not exploit.

For a computation one can use any candidate feasible rational solution of (II.18.3); an optimality proof is unnecessary when its objective is already below one. This separates discovering a certificate from checking one. The supplied script checks the exact primal–dual profile identity for many finite instances, but does not include a general LP solver.

II.19 Further proved extensions

Source 07, delivered Section 8.

II.19.1 Recovering labels from collision probabilities

So far the labels $\psi(d)$ were given. They can instead be selected from a concentrated directional difference distribution. For $h \in H$, define

$$p_d(h) = \sum_{\substack{g \in S \\ g+d \in S}} \mu(g) \frac{\#\{(x, y) \in X_g \times X_{g+d} : \phi(y) - \phi(x) = h\}}{R_g R_{g+d}}. \quad (\text{II.19.1})$$

This distribution has finite support even if H is infinite, and total mass $s_d = 1 - b_d$.

Proposition II.19.1 (Collision-to-mode exactification). *Assume $s_d > 0$ and put*

$$C_d = \sum_{h \in H} \left(\frac{p_d(h)}{s_d} \right)^2.$$

Choose $\psi(d)$ to maximize $p_d(h)$. Its directional defect satisfies

$$\delta_d = 1 - \max_h p_d(h) \leq 1 - s_d C_d. \quad (\text{II.19.2})$$

In particular, if $C_d \geq 1 - \varepsilon_d$, then

$$\delta_d \leq b_d + (1 - b_d)\varepsilon_d.$$

Substituting these budgets into Theorem II.15.1 or Corollary II.15.3 gives the corresponding exact relations among the selected modes.

Proof. For $q_h = p_d(h)/s_d$, the values q_h are a probability distribution. Hence

$$C_d = \sum_h q_h^2 \leq (\max_h q_h) \sum_h q_h = \max_h q_h.$$

Multiplying by s_d and using the definition of the defect proves the claim. If a selected label has defect below $1/2$, it is the unique label with that property, since two disjoint success events cannot both have probability greater than $1/2$. $\square$

[write] **6 October 2026; sources 11–14.** Source 12’s Lemma II.36.1 gives the sharp form of this collision-to-mode step, mode at least $1 - \chi(\eta)$ with $\chi(\eta) = (1 - \sqrt{1 - 2\eta})/2$; source 11’s Lemma II.28.1 uses the coarser $\sum p^2 \leq \max p$.

The collision experiment is two independent samples from the normalized internal difference law. It is not automatically the same as a raw additive-energy statistic, which weights pairs of fibers differently. An application starting from energy must separately prove the necessary normalization or comparison. This proposition does not bypass that step.

II.19.2 Weighted conditional fiber laws

The finite uniform model is not essential to the main theorem. More generally, put any probability law π_g on each nonempty usable fiber, and any positive probability law μ on the finite support S . Define $F_d(g)$ by sampling independently from π_g and π_{g+d} on internal edges, and again set it equal to one on exits. Define δ_d by its μ -mean and keep the same transport profiles.

Proposition II.19.2 (Weighted version). *Theorems II.15.1 and II.18.3 remain valid for these arbitrary conditional laws. The fiber-size estimates of Lemma II.14.4 are replaced, where needed, by direct bounds on $\mu(g - s)/\mu(g)$.*

Proof. Choose the initial index with the stated law and every formal fiber point with its prescribed conditional law, still sharing the endpoints and keeping distinct formal internal positions independent. The conditional edge-failure probability is the newly defined $F_d(g)$. The proof of (II.15.4) and all transport optimizations depend only on its range and mean, not on uniformity of the fiber samples. $\square$

II.19.3 Nonabelian source and target groups

Theorem II.19.3 (Ordered word relations). *Let G be a finite group and H any group, now written multiplicatively. Use right steps $g \mapsto gd$, and test $\phi(x)^{-1}\phi(y) = \psi(d)$ on an internal edge. Define the transport profile with coefficient $\mu(gs^{-1})$ for prefix s . For two nonempty words with the same ordered product in G , a total profile budget below one forces equality of their ordered image products in H .*

Proof. For a prefix s , the equation $g = rs$ gives starting index $r = gs^{-1}$, so the edge marginal has precisely the claimed coefficient. Choose shared initial and terminal samples exactly as in Theorem II.15.1. On the event that all edge equations hold,

$$\prod_{j=1}^m \psi(d_j) = [\phi(x_0)^{-1}\phi(x_1)][\phi(x_1)^{-1}\phi(x_2)] \cdots [\phi(x_{m-1})^{-1}\phi(x_m)] = \phi(x_0)^{-1}\phi(x_m),$$

with products in their original order. Both paths have the same endpoints. The identical union bound now proves the result without any commutativity assumption. $\square$

This is a relation-wise, direction-wise statement. It should again be distinguished from a theorem deriving closeness to a global homomorphism from one averaged multiplication test, as in [27].

II.20 An explicit weighted replacement for the energy extraction step

Source 10, delivered Section 7.

The rather large powers in Proposition 7.3 of the source paper are not intrinsic to the energy extraction step. Two intermediate thresholding operations can be removed by retaining the representation multiplicities as weights. We give all constants, including the finite-size correction caused by diagonal pairs. This is a self-contained quantitative replacement for that proposition, not a claim to improve the modern Balog–Szemerédi–Gowers theorem.

For a nonempty finite subset A of an abelian group, write

$$m = |A|, \quad r_A(d) = |\{(a, b) \in A^2 : a - b = d\}|, \quad E(A) = \sum_d r_A(d)^2.$$

The ambient group need not be finite or torsion-free.

Theorem II.20.1 (An explicit weighted energy extraction). *Suppose $0 < c \leq 1$ and $E(A) \geq cm^3$. Then there is a nonempty subset $A_* \subseteq A$ such that*

$$|A_*| \geq \frac{c}{4}m, \quad |A_* - A_*| \leq 2^{10}c^{-5}m.$$

Proof. If $m \leq 32c^{-2}$, take $A_* = A$. Indeed, $|A - A| \leq m^2 \leq 32c^{-2}m \leq 2^{10}c^{-5}m$. We henceforth assume $m > 32c^{-2}$.

Define the symmetric matrix of weights

$$w(a, b) = \frac{r_A(a - b)}{m} \in [0, 1] \quad (a, b \in A)$$

and its weighted codegrees

$$T(a, b) = \sum_{v \in A} w(v, a)w(v, b).$$

Here diagonal pairs are retained throughout. The normalization was chosen so that

$$\sum_{a, b \in A} w(a, b) = \frac{E(A)}{m} \geq cm^2.$$

Call a pair (a, b) bad if $T(a, b) < c^2m/16$.

Choose v uniformly from A and, conditional on v , form a random subset $K \subseteq A$ by including each a independently with probability $w(v, a)$. If $d(v) = \sum_a w(v, a)$, then

$$\mathbb{E}(|K|^2 \mid v) = d(v)^2 + \sum_a w(v, a)(1 - w(v, a)) \geq d(v)^2.$$

Cauchy–Schwarz therefore gives

$$\mathbb{E}|K|^2 \geq \frac{1}{m} \sum_v d(v)^2 \geq c^2m^2.$$

Let Z be the number of ordered bad pairs in K^2 . For $a \neq b$, the probability that $a, b \in K$ is $T(a, b)/m$, hence is less than $c^2/16$ for a bad pair. There are at most m diagonal pairs, and each contributes at most one to the expectation. Consequently

$$\mathbb{E}Z \leq \frac{c^2m^2}{16} + m.$$

It follows that

$$\mathbb{E}(|K|^2 - 8Z) \geq \frac{c^2 m^2}{2} - 8m > \frac{c^2 m^2}{4}.$$

Fix one outcome with $|K|^2 - 8Z \geq c^2 m^2/4$. This implies

$$|K| \geq cm/2, \quad Z \leq |K|^2/8.$$

Let A_* consist of those $a \in K$ for which at most $|K|/4$ elements $b \in K$ make (a, b) bad. Counting bad pairs by their first coordinate shows that $|A_*| \geq |K|/2 \geq cm/4$. Moreover, for every $a, b \in A_*$ there are at least $|K|/2$ elements $z \in K$ for which

$$T(a, z) \geq c^2 m/16, \quad T(b, z) \geq c^2 m/16.$$

Let $R(t)$ count the representations of t in the form $(p - q) - (r - s)$ with $p, q, r, s \in A$. For fixed $a, z \in A$,

$$R(a - z) \geq \sum_{u \in A} r_A(a - u) r_A(z - u) = m^2 T(a, z).$$

For completeness, the inequality involves no multiplicity error: each choice of u , a representation $p - q = a - u$, and a representation $r - s = z - u$ supplies a quadruple representing $a - z$; that quadruple recovers u from $u = a - (p - q)$.

For each of the above good choices of z , there are at least $c^4 m^6/2^8$ pairs of quadruples representing $a - z$ and $b - z$. Combining them gives representations of $a - b$ by ordered eight-tuples in A^8 . Representations coming from different z are disjoint: the first quadruple determines $a - z$ and hence determines z . Thus every $a - b$ with $a, b \in A_*$ has at least

$$\frac{|K|}{2} \frac{c^4 m^6}{2^8} \geq \frac{c^5 m^7}{2^{10}}$$

representations of this form. Choose one pair of endpoints (a, b) for every distinct member of $A_* - A_*$. An ordered eight-tuple has a unique value of the displayed alternating sum, so tuples counted for different differences are disjoint. There are only m^8 eight-tuples in total. Dividing proves $|A_* - A_*| \leq 2^{10} c^{-5} m$. $\square$

What is improved, and what is already known. The source paper gives $|A_*| \geq 2^{-20} c^{12} m$ and $|A_* - A_*| \leq 2^{38} c^{-24} m$. Theorem II.20.1 improves both quantitative guarantees and applies to arbitrary abelian groups. The mechanism is familiar: weighted dependent random choice, followed by counting many representations of each surviving difference. It should not be described as a new best theorem in additive combinatorics. For example, Reiher and Schoen prove that for every $\varepsilon \in (0, 1/2)$ one can obtain

$$|A'| \geq (1 - \varepsilon) c^{1/2} m, \quad |A' - A'| \leq 2^{33} \varepsilon^{-9} c^{-4} |A'|.$$

Their theorem is a substantially stronger modern benchmark [12, Theorem 1.2]. The advantage of the present replacement is a short proof with small, completely explicit constants that can be inserted directly into the source paper's chain of lemmas.

[write] **6 October 2026; sources 06–10.** Source 03's Theorem II.6.3 is the other elementary replacement, with $(c/8, 2^{19} c^{-9})$; this theorem is stronger in both conclusions for $0 < c \leq 1$ (note after source 03's theorem). Both imply the formal `proposition_7_3` bounds (`Sections06_07.lean`, line 135; exact companion `proposition_7_3_holds`); neither is formalized.

A parameterized version for later optimization. The same proof can retain its choices. Let $0 < \beta < \eta < 1/2$, $0 < \theta < \beta$, and $0 < \sigma < 1$ satisfy $\sigma^2 < 1 - \theta/\beta$. Put

$$M_0 = \frac{1}{\beta(1 - \theta/\beta - \sigma^2)c^2}.$$

If $m > M_0$, the proof with bad threshold $\theta c^2 m$ and potential $|K|^2 - \beta^{-1}Z$ gives

$$|A_*| \geq (1 - \beta/\eta)\sigma cm, \quad |A_* - A_*| \leq \frac{c^{-5}m}{(1 - 2\eta)\sigma\theta^2}.$$

For $m \leq M_0$ one can simply take $A_* = A$ and use $|A - A| \leq M_0 m$. The choice $(\beta, \eta, \theta, \sigma) = (1/8, 1/4, 1/16, 1/2)$ gives the theorem. The explicit m term in the random-pair calculation is essential: independence of the two membership events is available only off the diagonal. Omitting that term would leave a genuine gap for small sets.

II.21 A sharper obstruction-fibre form of the Freiman extraction

Source 10, delivered Section 8.

We next refine Lemma 7.5 of the source. This stage does not require the domain to be a cyclic group. What matters is that the *codomain* is the field of prime order.

Theorem II.21.1 (Extraction controlled by the vertical obstruction fibre). *Let G be an abelian group, let p be prime, let $B \subseteq G$ be finite and nonempty, and let $\phi : B \rightarrow \mathbb{F}_p$. Fix an integer $k \geq 1$, write $\Gamma = \{(x, \phi(x)) : x \in B\}$, and define*

$$H_k = \{y \in \mathbb{F}_p : (0, y) \in k\Gamma - k\Gamma\}, \quad h_k = |H_k|.$$

Then ϕ restricts to a Freiman homomorphism of order k on a subset $B' \subseteq B$ of size at least

$$|B'| \geq \frac{|B|}{k(h_k - 1) + 1}.$$

If, in addition, $|\Gamma - \Gamma| \leq C|\Gamma|$, then

$$h_k \leq C^{2k+1}, \quad |B'| \geq \frac{|B|}{kC^{2k+1}}.$$

Proof. The set H_k is symmetric and contains zero. If $h_k = 1$, then every relation between two sums of k elements of B already has zero vertical defect, so ϕ itself is a Freiman homomorphism of order k . Suppose henceforth that $h_k \geq 2$, and set

$$a = k(h_k - 1), \quad L = \left\lfloor \frac{p-2}{a} \right\rfloor.$$

For each integer j with $1 \leq j \leq kL$ and each $y \in H_k \setminus \{0\}$, the equation $jd = y$ excludes at most one $d \in \mathbb{F}_p^\times$. Indeed $kL \leq p-2$, so j is invertible. There are at most $kL(h_k - 1) = aL \leq p-2$ excluded nonzero choices of d , and there are $p-1$ nonzero choices available. Choose an allowed d . Symmetry of H_k ensures

$$jd \notin H_k \quad \text{whenever } 0 < |j| \leq kL.$$

For every $b \in \mathbb{F}_p$, consider

$$B_b = \{x \in B : \phi(x) \in b + d\{0, 1, \dots, L\}\}.$$

Since $L < p$ and $d \neq 0$, each $x \in B$ is counted in exactly $L + 1$ of these sets. Some b thus satisfies $|B_b| \geq (L + 1)|B|/p$. We claim

$$\frac{L + 1}{p} \geq \frac{1}{a + 1}.$$

If $a \geq p - 1$, then $L = 0$ and the assertion is immediate. If $a \leq p - 1$, Euclidean division of $p - 2$ by a gives $L + 1 \geq (p - 1)/a$, while $(p - 1)/a \geq p/(a + 1)$ because $p - 1 \geq a$.

Take $B' = B_b$. For each $x \in B'$ there is a unique integer $j(x) \in [0, L]$ such that $\phi(x) = b + dj(x)$. If $x_1 + \dots + x_k = x'_1 + \dots + x'_k$, then

$$d \left(\sum_{i=1}^k j(x_i) - \sum_{i=1}^k j(x'_i) \right) \in H_k.$$

The integer in parentheses lies in $[-kL, kL]$, and the chosen step d forces it to vanish. Hence the corresponding sums of ϕ values are equal, proving the Freiman property and the first bound.

For the small-doubling assertion, the map

$$H_k \times \Gamma \longrightarrow (k + 1)\Gamma - k\Gamma, \quad (y, \gamma) \longmapsto (0, y) + \gamma$$

is injective. Its first coordinate recovers γ because Γ is a graph, and then its second coordinate recovers y . The standard Plünnecke–Ruzsa inequality gives

$$|(k + 1)\Gamma - k\Gamma| \leq C^{2k+1}|\Gamma|.$$

For a small-difference hypothesis this follows by applying the sum-difference form of that inequality to the pair $-\Gamma, \Gamma$; see [36]. Consequently $h_k \leq C^{2k+1}$. Finally, $k(h_k - 1) + 1 \leq kC^{2k+1}$ for $k \geq 1$. $\square$

[write] **6 October 2026; sources 21–27.** Source 24’s Theorem II.79.3 is a second route with the prefactor $|B|/(kQ_k)$, $Q_k = h_k$, never better than the one here; its Theorem II.79.5 treats codomains $\mathbb{F}_p^d$ by linear hashing.

[write] **6 October 2026; sources 28–33.** Source 31’s Theorem II.116.1 is this theorem with arbitrary nonnegative weights and the progression length $\lfloor p/(kq) \rfloor$, never longer than the one in the proof here (note there).

Why the fibre is the appropriate invariant. The set H_k records exactly the possible failures to respect length- k additive relations. The global size of a high sum-difference set is used only to bound h_k when no more specific information is available. In examples where the vertical defects occupy few values, the first bound in Theorem II.21.1 can be much stronger than its small-doubling corollary. If $H_k = \{0\}$ it correctly returns all of B , independently of the doubling constant. The dilation proof requires a prime codomain; the displayed argument does not extend unchanged to arbitrary composite cyclic groups.

Corollary II.21.2 (An explicit replacement for Corollary 7.6). *Let B_0 be a nonempty finite subset of an abelian group, let $\phi : B_0 \rightarrow \mathbb{F}_p$, and suppose the graph of ϕ has at least $\gamma|B_0|^3$ additive*

quadruples, where $0 < \gamma \leq 1$. For every integer $k \geq 1$ there is $B \subseteq B_0$ on which ϕ is a Freiman homomorphism of order k and

$$|B| \geq \frac{2^{-24k-14}}{k} \gamma^{12k+7} |B_0|.$$

In particular, for order eight,

$$|B| \geq 2^{-209} \gamma^{103} |B_0|.$$

Proof. Apply Theorem II.20.1 to the graph, obtaining a subgraph Γ_1 with

$$|\Gamma_1| \geq \frac{\gamma}{4} |B_0|, \quad |\Gamma_1 - \Gamma_1| \leq 2^{10} \gamma^{-5} |B_0| \leq 2^{12} \gamma^{-6} |\Gamma_1|.$$

Set $C = 2^{12} \gamma^{-6}$ and use Theorem II.21.1. The retained proportion is at least

$$\frac{\gamma/4}{k(2^{12} \gamma^{-6})^{2k+1}} = \frac{2^{-24k-14}}{k} \gamma^{12k+7}.$$

When $k = 8$, $24k + 14 + \log_2 k = 209$ and $12k + 7 = 103$. □

[write] 6 October 2026; sources 06–10. *Checked at the second write.* The retained proportion is $(\gamma/4)/(k(2^{12} \gamma^{-6})^{2k+1}) = 2^{-2-12(2k+1)} k^{-1} \gamma^{12k+7}$, which at $k = 8$ is $2^{-209} \gamma^{103}$. *Against the formal statements* (HEAD 764740f07): `corollary_7_6` (Sections06_07.lean, line 197) assumes N prime, $|B_0| = \alpha N$ and $\gamma(\alpha N)^3 \leq$ the number of additive quadruples of the graph, and concludes a Freiman homomorphism of order eight on B with $|B| \geq 2^{-1882} \gamma^{1164} \alpha N$. The hypothesis forces $\gamma \leq 1$ (a graph on B_0 has at most $|B_0|^3$ additive quadruples), and then $2^{-209} \gamma^{103} \geq 2^{-1882} \gamma^{1164}$, so this corollary implies the formal conclusion with $\mathbb{F}_p = \mathbb{Z}/N\mathbb{Z}$. Likewise Theorem II.21.1 implies `lemma_7_5`'s $|B|/(8kC^{4k})$, since $C \geq 1$ when $|\Gamma - \Gamma| \leq C|\Gamma|$ and $B \neq \emptyset$ (for $B = \emptyset$ the formal conclusion holds with $B' = \emptyset$) (corrected after the independent check of 6 October 2026; the first wording omitted the empty case). Neither is formalized (exact companions `corollary_7_6_holds`, `lemma_7_5_holds` prove the printed constants).

[write] 6 October 2026; sources 15–20. This partly answers source 17's Question 7 (exactness from quadruple energy with polynomial loss; item II.36).

[write] 6 October 2026; sources 21–27. Source 24's Theorem F.29.1 gives $2^{-(34k+20)} \gamma^{12k+7}/k$, weaker than this corollary by 2^{10k+6} ; its Corollary II.80.2 is identical to the modern-input alternative below.

[write] 6 October 2026; sources 28–33. Source 31's elementary route gives $2^{-363} c^{171}$ at order eight with weights (Theorem II.116.2, $K = 1$), weaker than $2^{-209} \gamma^{103}$ here; its Reiher–Schoen bound (II.116.5) is the modern-input alternative below.

A modern-input alternative. If the Reiher–Schoen theorem [12] is admitted as an external input instead of the self-contained energy extraction, the same fibre argument gives

$$|B| \geq \frac{(1-\varepsilon)\varepsilon^{9(2k+1)}}{k 2^{33(2k+1)}} \gamma^{8k+9/2} |B_0| \quad (0 < \varepsilon < 1/2).$$

For $k = 8$ the power of γ is $137/2$ rather than 103. This is a corollary of a published result plus the preceding extraction calculation, and must be labelled as such. The self-contained and modern-input versions have different useful roles: the former has a short complete proof here, while the latter records a better presently available exponent. Neither by itself yields a new overall bound for Szemerédi's theorem; the later structural and iterative stages still have to be checked with their own quantified hypotheses.

Further questions arising specifically from these refinements. The obstruction fibres H_k may permit data-sensitive structure theorems stronger than replacing them immediately by C^{2k+1} . Natural precise projects are to bound $|H_k|$ from a graph's vertical energy profile, to study its growth with k , and to determine when an order- k homomorphism extends with small additional loss to order $k + 1$. The next subsection extends the prime-codomain dilation argument to a general finite abelian codomain by retaining the exact number of solutions of $jd = y$. For formalization, the two finite counting injections in the energy proof and the injection $H_k \times \Gamma \rightarrow (k + 1)\Gamma - k\Gamma$ are particularly suitable intermediate lemmas.

II.21.1 A torsion-sensitive extraction for general finite codomains

The prime-codomain dilation argument admits an explicit extension when the sizes of the multiplication fibres are retained. The resulting criterion records how the vertical defects meet the images jV .

Theorem II.21.3 (Extraction from a torsion and defect budget). *Let G be an abelian group, let $B \subseteq G$ be finite and nonempty, let V be a finite abelian group of order $n \geq 2$, and let $\phi : B \rightarrow V$. Fix an integer $k \geq 1$, and put*

$$\Gamma = \{(x, \phi(x)) : x \in B\}, \quad F = \{y \in V : (0, y) \in k\Gamma - k\Gamma\}.$$

For an integer $L \geq 0$ define

$$\mathcal{B}_{k,V,F}(L) = \sum_{j=1}^{kL} (|V[j]| |F \cap jV| - 1), \quad V[j] = \{v \in V : jv = 0\}.$$

If

$$\mathcal{B}_{k,V,F}(L) \leq n - 2, \tag{II.21.1}$$

then ϕ restricts to a Freiman homomorphism of order k on a subset $B' \subseteq B$ of size at least

$$|B'| \geq \frac{L+1}{n} |B|. \tag{II.21.2}$$

The subset can be chosen so that its image lies in $a + d\{0, \dots, L\}$ for some $a \in V$ and some $d \neq 0$. If $L > 0$, this step has order greater than kL . If $F = \{0\}$, the original map ϕ is already a Freiman homomorphism of order k on all of B , independently of the budget.

Proof. The defect set F is symmetric and contains zero. For $j \geq 1$ let

$$E_j = \{d \in V \setminus \{0\} : jd \in F\}.$$

Multiplication by j has image jV , and every nonempty fibre has $|V[j]|$ elements. Because $0 \in F$, the zero element is one of the preimages being counted. Removing it gives the exact identity

$$|E_j| = |V[j]| |F \cap jV| - 1. \tag{II.21.3}$$

The union bound and (II.21.1) leave at least one nonzero d outside $\bigcup_{j=1}^{kL} E_j$. It satisfies

$$jd \notin F \quad \text{for } 1 \leq j \leq kL.$$

By symmetry of F , the same holds for $-kL \leq j \leq -1$. When $L > 0$, the inclusion $0 \in F$ also implies that the order of d is greater than kL . Hence the $L + 1$ elements of $P = d\{0, \dots, L\}$ are distinct; this is immediate when $L = 0$.

For each $a \in V$ let

$$B_a = \{x \in B : \phi(x) \in a + P\}.$$

Every $x \in B$ occurs in exactly $L + 1$ of these sets, since the possible translations are $a = \phi(x) - jd$ with $0 \leq j \leq L$. Thus

$$\sum_{a \in V} |B_a| = (L + 1)|B|.$$

Choose a translate attaining at least this average, and put $B' = B_a$. For every $x \in B'$ there is a unique integer $t(x) \in [0, L]$ with $\phi(x) = a + dt(x)$.

Suppose $x_1 + \dots + x_k = x'_1 + \dots + x'_k$ with all the points in B' . Its vertical defect belongs to F and equals

$$dJ, \quad J = \sum_{i=1}^k t(x_i) - \sum_{i=1}^k t(x'_i), \quad |J| \leq kL.$$

The choice of d rules out every nonzero such integer J . Therefore $J = 0$, and the sums of the corresponding ϕ values agree. This proves the Freiman property and (II.21.2). Finally, if $F = \{0\}$, every length- k relation in the original domain has zero vertical defect by definition, proving the final assertion. $\square$

[write] **6 October 2026; sources 21–27.** On $V = \mathbb{F}_p^d$ source 24's Theorem II.79.5 is strictly better whenever $Q_k \leq p^{d-1}$ and $k \geq 2$, because a multiple of p costs $p^d - 1$ in the budget here (proved in the note after that theorem).

Corollary II.21.4 (Recovery of the prime-codomain estimate). *Under the hypotheses above, let $V = \mathbb{F}_p$ and $h = |F|$. There is a subset on which ϕ is a Freiman homomorphism of order k and whose cardinality is at least*

$$\frac{|B|}{k(h - 1) + 1}.$$

Proof. If $h = 1$, use all of B . Otherwise set

$$a = k(h - 1), \quad L = \left\lfloor \frac{p - 2}{a} \right\rfloor.$$

Then $kL \leq p - 2$, so multiplication by each $1 \leq j \leq kL$ is invertible. Every summand of the budget is $h - 1$, and

$$\mathcal{B}_{k, \mathbb{F}_p, F}(L) = kL(h - 1) \leq p - 2.$$

The theorem gives $(L + 1)|B|/p$. The integer inequality $p - 2 < a(L + 1)$ implies

$$p \leq a(L + 1) + 1 \leq (a + 1)(L + 1),$$

which proves the claimed bound, including the case $L = 0$. $\square$

Scope of the criterion. The budget always allows $L = 0$, giving the constant-fibre bound $|B|/|V|$. Positive L requires a step of order greater than kL . For example, if $\exp(V) \leq k$, no positive L satisfies the budget: the summand with $j = \exp(V)$ is already $n - 1$. The separate case $F = \{0\}$ still returns the whole domain. Thus the theorem is a data-dependent extension of the dilation argument; it does not assert a uniform large-subset guarantee for every torsion profile. Sharper estimates for unions of the sets E_j , rather than their sum, are a natural way to improve this criterion.

II.22 Source 11’s chapter: affine rigidity at the Section 6 endpoint [write]

Added in the third write, 6 October 2026.

The next ten sections are source 11, a manuscript on one interface of Gowers’s Section 6: after Proposition 6.1 supplies many respected additive quadruples and Lemma 6.2 identifies them with additive quadruples of the graph of the frequency-selection map, the observation before Lemma 6.2 says that *full* graph energy forces an affine map. Source 11 quantifies this endpoint. *Logically the chapter follows source 02’s Sections II.1–II.4*, which strengthen the energy lower bound of Proposition 6.1 itself; it is printed here, at the end of Part II, because no number printed earlier may move. Its introduction is Section F.16; its verification, formalization plan, questions, conclusion and reproduction notes are in Part V (Sections V.46–V.49).

Reading conventions. $\mathcal{E}(f)$ counts ordered additive quadruples respected by $f : G \rightarrow W$, $\varepsilon(f) = 1 - \mathcal{E}(f)/N^3$, $d(f)$ is the proportion of points where f differs from the nearest affine map, $\kappa = 1$ exactly when W has an element of order two, and $\Psi_\kappa(t) = 4t - (10 + 2\kappa)t^2 + (6 + 2\kappa)t^3$; the table at the head of this Part lists the clashes.

What is proved, in one paragraph. For errors on a set of density t below $1/7$ (odd targets) or $1/3$ (targets with an involution) the defect is at least $\Psi_\kappa(t)$, with equality exactly for one error value on a subgroup coset (Theorem F.16.1); if $\varepsilon(f) < 1/9$ the nearest affine map is unique and $d(f) \leq D_\kappa(\varepsilon(f))$ (Theorem F.16.2); for $|G| \geq 34$ the largest non-affine graph energy is $N^3 - 4N^2 + (10 + 2\kappa)N - (6 + 2\kappa)$, attained exactly by one-point modifications, and for maps between vector spaces over the same prime field without the cutoff (Theorem F.16.3). The radii $1/7$, $1/3$ and the cutoff 34 are sufficient, not claimed optimal (the source’s non-claims); Remark W2 shows that the cutoff cannot simply be dropped for general groups. The correction step is the classical Blum–Luby–Rubinfeld majority decoding, credited by the source.

Relation to the formal statements (checked at HEAD e1db00e99). `proposition_6_1` (Section ns06_07.lean, line 37; exact companion `proposition_6_1_holds`, Proofs05_10.lean, line 505) and `lemma_6_2` (line 47; `lemma_6_2_holds`, Proofs05_10.lean, line 589, with `Finset.addEnergy` of `functionGraph`) are the interface. The exact endpoint itself is a remark of the paper, not a numbered statement or a Lean declaration. Nothing in this chapter is formalized; the source’s plan is in Part V and in 11-affine-rigidity-formalization_plan.md.

II.23 Graph energy, derivatives, and the exact endpoint

Source 11, delivered Section 2.

Lemma II.23.1 (Derivative collision identity). *Put $\Delta_h f(x) = f(x + h) - f(x)$ and*

$$r_h(w) = |\{x \in G : \Delta_h f(x) = w\}|.$$

Then

$$\mathcal{E}(f) = \sum_{h \in G} \sum_{w \in W} r_h(w)^2. \quad (\text{II.23.1})$$

Moreover, adding or subtracting an affine map from f does not change $\mathcal{E}(f)$.

Proof. Pairs x, y with $\Delta_h f(x) = \Delta_h f(y)$ correspond bijectively to respected quadruples $(x + h, y, x, y + h)$. Conversely, from any respected quadruple one recovers $h = x_1 - x_3 = x_4 - x_2$, $x = x_3$, and $y = x_2$. For an affine map $\ell(x) = L(x) + b$, the derivative $\Delta_h \ell(x) = L(h)$ is independent of x . Subtracting it just permutes the labels of each histogram r_h . $\square$

Lemma II.23.2 (Full energy). *One has $\mathcal{E}(f) \leq N^3$, with equality if and only if f is affine.*

Proof. For each h , the nonnegative integers $r_h(w)$ sum to N , so their squares sum to at most N^2 . Equality holds precisely when the derivative is constant. If all derivatives are constant, let their values be $L(h)$. The identity

$$\Delta_{h+k} f(x) = \Delta_h f(x + k) + \Delta_k f(x) \quad (\text{II.23.2})$$

implies $L(h + k) = L(h) + L(k)$. Also $f(x) = L(x) + f(0)$. The converse is immediate. $\square$

In the derivative-imbalance normalization used in [50], if $M = |W|$, the corresponding quantity is

$$NB_f = \mathcal{E}(f) - N^2 - (N - 1)N^2/M.$$

The ambiguity statistic is

$$\sum_{h \neq 0} \sum_w \binom{r_h(w)}{2} = \frac{\mathcal{E}(f) - 2N^2 + N}{2}.$$

These formulas allow the results here to be translated into those normalizations without confusing componentwise Fourier nonlinearity with the distance from a vector-valued map to a single affine map.

II.23.1 A basic fact about error supports

For a subset $S \subseteq G$ let

$$E(S) = \#\{(x_1, x_2, x_3, x_4) \in S^4 : x_1 + x_2 = x_3 + x_4\}.$$

This is set energy, as distinct from the graph energy $\mathcal{E}(f)$.

Lemma II.23.3 (Maximum set energy and its equality case). *If $|S| = s$, then $E(S) \leq s^3$. If $S \neq \emptyset$, equality holds if and only if S is a coset of a subgroup of G .*

Proof. Write $c_h = |S \cap (S - h)|$. Then $E(S) = \sum_h c_h^2$ and $\sum_h c_h = s^2$. Since $0 \leq c_h \leq s$, $E(S) \leq s \sum_h c_h = s^3$. Equality forces each c_h to be either 0 or s . Let $T = \{h : S + h = S\}$, a subgroup. Fix $a \in S$. For every $x \in S$, $c_{x-a} > 0$, so $x - a \in T$. Thus $S \subseteq a + T$, while $a + T \subseteq S$ follows from translation invariance. A subgroup coset plainly attains equality. $\square$

II.24 Exact counting for a one-valued error

Source 11, delivered Section 3.

Let $S \subseteq G$, $|S| = s$, and $b \neq 0$ in W . Consider $g = b\mathbf{1}_S$. Every additive quadruple has four positions, with signs $+, +, -, -$. Let J denote the set of positions whose coordinates belong to S .

Any choice of three coordinates uniquely determines the fourth, and any choice of two coordinates has exactly N completions. This remains true when coordinates coincide or G has torsion. Consequently, fixing three specified positions in S gives s^3 quadruples, fixing two gives s^2N , and fixing one gives sN^2 .

Lemma II.24.1 (Support-pattern counts). *The numbers of additive quadruples with exactly one, exactly three, or exactly two same-side positions in S are, respectively,*

$$Q_1 = 4sN^2 - 12s^2N + 12s^3 - 4E(S), \quad (\text{II.24.1})$$

$$Q_3 = 4s^3 - 4E(S), \quad (\text{II.24.2})$$

$$Q_{2,\text{same}} = 2s^2N - 4s^3 + 2E(S). \quad (\text{II.24.3})$$

Proof. For one prescribed position in S and the remaining three outside S , inclusion–exclusion gives $sN^2 - 3s^2N + 3s^3 - E(S)$; sum over the four positions. For a prescribed triple of positions in S and the remaining position outside, the answer is $s^3 - E(S)$; again there are four choices. The two same-side pairs are $\{1, 2\}$ and $\{3, 4\}$. For either pair, requiring the other two positions to lie outside S gives $s^2N - 2s^3 + E(S)$. $\square$

Proposition II.24.2 (Exact one-valued defect). *If $2b \neq 0$, then*

$$N^3 - \mathcal{E}(b\mathbf{1}_S) = 4sN^2 - 10s^2N + 12s^3 - 6E(S). \quad (\text{II.24.4})$$

If $2b = 0$, then

$$N^3 - \mathcal{E}(b\mathbf{1}_S) = 4sN^2 - 12s^2N + 16s^3 - 8E(S). \quad (\text{II.24.5})$$

Proof. When $2b \neq 0$, the rejected support patterns have exactly one position, exactly three positions, or exactly two positions on the same side. Patterns with two opposite-side positions have equal contributions on both sides, and patterns with zero or four positions also pass. Thus the defect is $Q_1 + Q_3 + Q_{2,\text{same}}$.

When $2b = 0$, the same-side two-position patterns pass as well. The rejected patterns are exactly those with an odd number of positions in S , giving $Q_1 + Q_3$. $\square$

Write $t = s/N$ and $e_S = E(S)/N^3$. The two formulas become

$$\varepsilon(b\mathbf{1}_S) = \Psi_0(t) + 6(t^3 - e_S), \quad 2b \neq 0, \quad (\text{II.24.6})$$

$$\varepsilon(b\mathbf{1}_S) = \Psi_1(t) + 8(t^3 - e_S), \quad 2b = 0. \quad (\text{II.24.7})$$

The set-energy gap is therefore visible without any asymptotic notation.

Example II.24.3 (One corrupted point). Take $S = \{x_0\}$. Then $E(S) = 1$, and

$$\mathcal{E}(b\mathbf{1}_{\{x_0\}}) = \begin{cases} N^3 - 4N^2 + 10N - 6, & 2b \neq 0, \\ N^3 - 4N^2 + 12N - 8, & 2b = 0. \end{cases}$$

For $h \neq 0$, the derivative histogram has entries $N - 2, 1, 1$ in the first case and $N - 2, 2$ in the second. This gives an independent check of the counting formulas. The distinction depends on the target, not merely on the parity of the domain size.

II.25 Replacing many error values by one: odd targets

Source 11, delivered Section 4.

Subtract an affine map and suppose now that g is zero outside S and nonzero on S . For $b \neq 0$ put

$$c_b = |\{x : g(x) = b\}|, \quad M = \sum_{b \neq 0} c_b^2, \quad A = s^2 - M, \quad C = \sum_{b \neq 0} c_b c_{-b}. \quad (\text{II.25.1})$$

The quantity A counts ordered pairs of errors with different nonzero values. It is zero exactly when all errors have one value.

Lemma II.25.1 (Two elementary color inequalities). *Let $T = s^3 - \sum_{b \neq 0} c_b^3$, the number of ordered non-monochromatic triples of error positions. Then*

$$T \leq \frac{3}{2} sA. \quad (\text{II.25.2})$$

For every target W , $C \leq M$. If W has no nonzero element of order two, one also has $C \leq A$.

Proof. A non-monochromatic triple has at least two unequal pairs among its three pairs of positions. Counting unequal pairs together with the third entry gives $3sA$, which proves (II.25.2). The inequality $C \leq M$ follows from

$$M - C = \frac{1}{2} \sum_{b \neq 0} (c_b - c_{-b})^2 \geq 0.$$

When $b \neq -b$ for every nonzero b , the ordered pairs counted by C form a subset of the unequal-value ordered pairs counted by A . $\square$

Proposition II.25.2 (Odd-target error-value comparison). *Suppose that W has no nonzero element of order two. Choose any $b_0 \neq 0$ and put $g_0 = b_0 \mathbf{1}_S$. Then*

$$\mathcal{E}(g_0) - \mathcal{E}(g) \geq 2(N - 7s)A. \quad (\text{II.25.3})$$

Proof. Separate additive quadruples according to the exact number of positions in S . With zero positions, both maps pass; with one position, both fail; with four positions, g_0 always passes and therefore cannot lose. Only the cases with two or three positions need attention.

For two opposite-side positions, g_0 passes, whereas g passes exactly when the two nonzero values agree. There are four opposite-side pairs. For one fixed pair and one of its A unequal-value assignments, the remaining coordinates have N possible completions. For either remaining coordinate, at most s completions put that coordinate in S , since it varies bijectively with the free parameter. Thus at least $N - 2s$ completions have exactly these two errors. The total gain is at least $4(N - 2s)A$.

For two same-side positions, g_0 fails. The map g can pass only when the values are negatives of one another. Ignoring the restriction on the other positions gives an upper bound $2NC$ for this loss, because there are two same-side pairs.

For exactly three error positions, g_0 fails. If all three nonzero values of g are equal, g fails as well: the signed sum is that value or its negative. Therefore a gain for g requires a non-monochromatic triple. There are four triples of positions, and dropping the restriction on the fourth position bounds this loss by $4T$.

Combining the three estimates and applying Lemma II.25.1,

$$\begin{aligned}\mathcal{E}(g_0) - \mathcal{E}(g) &\geq 4(N - 2s)A - 2NC - 4T \\ &\geq 4(N - 2s)A - 2NA - 6sA \\ &= 2(N - 7s)A.\end{aligned}$$

Repeated coordinates create no missing cases: the fixed-position parametrizations count them with their proper multiplicities. $\square$

Theorem II.25.3 (Odd-target deficit decomposition). *Let W have no nonzero element of order two, let $g = f - \ell$ for an affine ℓ , and use the notation above. Put $a = A/N^2$. Then*

$$\varepsilon(f) \geq \Psi_0(t) + 6(t^3 - e_S) + 2(1 - 7t)a. \quad (\text{II.25.4})$$

In particular, if $0 < t < 1/7$, the lower bound $\Psi_0(t)$ holds, and equality is equivalent to S being a subgroup coset and g having one nonzero value.

Proof. Affine invariance, Proposition II.24.2, and Proposition II.25.2 give (II.25.4). If $t < 1/7$, both remainder terms are nonnegative. Equality forces $A = 0$ and $E(S) = s^3$, giving the description by Lemma II.23.3. Conversely, that description attains equality by (II.24.6). $\square$

The comparison inequality is stronger than merely proving a cubic upper bound for graph energy. It gives two independently meaningful certificates of failure to be extremal: the support may have insufficient additive energy, and the errors may use too many different values.

II.26 Targets with two-torsion, including mixed targets

Source 11, delivered Section 5.

It is not necessary for the target to have characteristic two. In this section W may, for example, be $\mathbb{Z}/6\mathbb{Z}$; the hypothesis is only that it contains a nonzero involution b_0 . Keep M, A, C as in (II.25.1), and put

$$J = M - C \geq 0.$$

The extra statistic J measures the asymmetry between the multiplicities of b and $-b$.

Proposition II.26.1 (Error-value comparison with an involution). *If $s < N/3$ and $g_0 = b_0 \mathbf{1}_S$, with $2b_0 = 0$ and $b_0 \neq 0$, then*

$$\mathcal{E}(g_0) - \mathcal{E}(g) \geq 6(N - 3s)A + 2(N - 2s)J. \quad (\text{II.26.1})$$

Proof. For exactly two errors, g_0 always passes. For an opposite-side pair, g fails on A ordered assignments; for a same-side pair it fails on $s^2 - C = A + J$ assignments. Requiring the other two positions to be outside S leaves at least $N - 2s$ completions for each assignment. Hence the two-error gain is at least

$$(N - 2s)(4A + 2(A + J)).$$

For exactly three errors, the loss is bounded by $4T \leq 6sA$ just as before. With zero, one, or four errors, g_0 cannot lose. Subtraction proves the assertion. $\square$

Theorem II.26.2 (Two-torsion deficit decomposition). *Under these hypotheses, with $a = A/N^2$ and $j = J/N^2$,*

$$\varepsilon(f) \geq \Psi_1(t) + 8(t^3 - e_S) + 6(1 - 3t)a + 2(1 - 2t)j. \quad (\text{II.26.2})$$

If $0 < t < 1/3$, equality with $\Psi_1(t)$ holds precisely when the error is one involution-valued constant on a subgroup coset.

Proof. Combine (II.24.7) and (II.26.1). At equality all three remainder terms vanish. The condition $A = 0$ says that exactly one nonzero value b occurs. In this case $J = 0$ is equivalent to $b = -b$. The support condition again follows from Lemma II.23.3. The converse follows from the exact counting formula. $\square$

Remark II.26.3 (Why an odd-target assertion cannot simply ignore torsion). For a one-point error into a target containing an involution, Example II.24.3 exceeds the odd-target extremal energy by $2N - 2$. The odd-target cubic is therefore false in this setting, not just nonoptimal. The distinction has been incorporated at the level of the hypotheses rather than hidden in a proof convention.

II.27 Quantitative information about near-extremizers

Source 11, delivered Section 6.

The preceding theorems also describe how an almost extremizer can fail to be an exact extremizer. This does not require any compactness argument.

Corollary II.27.1 (Two separate defects in the odd case). *Assume $0 < t < 1/7$ and $\varepsilon(f) \leq \Psi_0(t) + \eta$. Then*

$$t^3 - e_S \leq \eta/6, \quad (\text{II.27.1})$$

$$\frac{A}{N^2} \leq \frac{\eta}{2(1 - 7t)}. \quad (\text{II.27.2})$$

There is a nonzero value b such that the proportion of domain points with nonzero error different from b is at most

$$\frac{\eta}{2t(1 - 7t)}. \quad (\text{II.27.3})$$

Proof. The first two claims follow by dropping the other nonnegative term in (II.25.4). If $m = \max_b c_b$ and $u = s - m$, then $M \leq sm$, so $A = s^2 - M \geq su$. Divide by N^2 and use (II.27.2). $\square$

For a target with an involution, (II.26.2) similarly gives

$$t^3 - e_S \leq \eta/8, \quad \frac{A}{N^2} \leq \frac{\eta}{6(1 - 3t)}, \quad \frac{J}{N^2} \leq \frac{\eta}{2(1 - 2t)}.$$

The minority-value estimate follows by dividing the second inequality by t . The third inequality is additional information not present in the odd-target case.

These estimates do not by themselves state that S is close in symmetric difference to a subgroup coset. They reduce that remaining stability issue to a precise high-energy question about S , with $E(S)/s^3 \geq 1 - \eta/(6t^3)$ in the odd case. That normalization matters: an additive error small compared with one need not be small compared with t^3 . A full relative stability statement would have to retain this scale.

II.28 Removing the assumption of proximity to an affine map

Source 11, delivered Section 7.

We now prove a correction lemma rather than taking it as a black box. The method is classical; its role here is to enter the range where the sharp cubics apply.

Lemma II.28.1 (Elementary correction of a nearly additive map). *Let $F : G \rightarrow W$ and suppose*

$$\mathbb{P}_{x,y}(F(x+y) \neq F(x) + F(y)) = \rho < 1/6.$$

Then there is a homomorphism $L : G \rightarrow W$ with

$$\mathbb{P}_x(F(x) \neq L(x)) \leq \frac{\rho}{1-2\rho}. \quad (\text{II.28.1})$$

Proof. Fix $h \in G$ and consider the random value $Z_h(y) = F(h+y) - F(y)$. For independent uniform y, z , both equations

$$F(h+y) + F(z) = F(h+y+z), \quad F(y) + F(h+z) = F(h+y+z)$$

hold except with probability at most 2ρ . Each tested pair is uniformly distributed on G^2 , even for fixed h . When both equations hold, $Z_h(y) = Z_h(z)$. Hence the collision probability of Z_h is at least $1 - 2\rho$.

For any probability distribution (p_w) , $\sum_w p_w^2 \leq \max_w p_w$. Thus Z_h has a unique modal value $L(h)$ of probability at least $1 - 2\rho > 1/2$. For fixed h, k , the three identities

$$Z_{h+k}(x) = L(h+k), \quad Z_h(x+k) = L(h), \quad Z_k(x) = L(k)$$

hold simultaneously for some x , since their total failure probability is at most $6\rho < 1$. The exact derivative identity $Z_{h+k}(x) = Z_h(x+k) + Z_k(x)$ then proves additivity of L .

If $F(h) \neq L(h)$, at least a proportion $1 - 2\rho$ of the values $Z_h(y)$ differ from $F(h)$. Averaging this observation over h gives

$$\rho = \mathbb{P}_{h,y}(Z_h(y) \neq F(h)) \geq (1 - 2\rho)\mathbb{P}_h(F(h) \neq L(h)),$$

which proves (II.28.1). □

[write] **6 October 2026; sources 11–14.** *Credit and comparison.* The majority step $\sum_w p_w^2 \leq \max_w p_w$ used here is the coarsest of three collision-to-mode bounds in this report: source 07's Proposition II.19.1 and source 12's sharp Lemma II.36.1, whose mode bound $1 - \chi(\eta)$, $\chi(\eta) = (1 - \sqrt{1 - 2\eta})/2 \leq \eta$, is optimal. Inserting source 12's lemma into this proof (with $\eta = 2\rho$) would give disagreement at most $\rho/(1 - \chi(2\rho))$ in (II.28.1); the consequence for the cutoff 34 of Theorem F.16.3 is recorded as a marked question, item II.21, because neither source writes it.

Lemma II.28.2 (Passing from graph energy to the correction test). *There is an anchor $a \in G$ such that $F_a(x) = f(a+x) - f(a)$ fails the test in Lemma II.28.1 with probability at most $\varepsilon(f)$.*

Proof. The average over a, x, y of the acceptance condition is the probability of

$$f(a+x) + f(a+y) = f(a) + f(a+x+y).$$

The parametrization $(a, x, y) \mapsto (a+x, a+y, a, a+x+y)$ is a bijection onto the ordered additive quadruples. Its acceptance probability is exactly $1 - \varepsilon(f)$. One anchor is at least as good as the average. □

Proof of Theorem F.16.2. Write $\varepsilon = \varepsilon(f)$. Applying Lemmas II.28.2 and II.28.1 gives an affine map with disagreement proportion at most

$$\frac{\varepsilon}{1 - 2\varepsilon} < \frac{1}{7} \quad (\varepsilon < 1/9).$$

In original coordinates this map is $x \mapsto L(x - a) + f(a)$. The closest affine map has disagreement proportion $d(f) < 1/7$ as well, so Theorem F.16.1 gives $\varepsilon \geq \Psi_\kappa(d(f))$. Monotonicity on $[0, 1/7]$ yields (F.16.5).

Two distinct affine maps agree either nowhere or on a coset of the kernel of a nonzero homomorphism. Such a coset has size at most $N/2$. Their mutual Hamming distance is therefore at least $N/2$. Two affine maps each at distance less than $N/7$ from f cannot be distinct. This proves uniqueness. Finally, equality in the inverse inequality is equality in the corresponding local profile, whose equality cases have already been classified. $\square$

II.28.1 A deterministic reconstruction algorithm

The proof gives an explicit, though not optimized, algorithm for a finite input table. Count the failed tests for each anchor a , select an anchor with failure at most the average, and compute the modal derivative value $L(h)$ for every h . Output $L(x - a) + f(a)$ and count its disagreements with f . With group operations available by tables, straightforward enumeration uses $O(N^3)$ elementary operations for the anchor search and $O(N^2)$ for the histograms, with sparse counters for target values. These operation counts describe this implementation, not a claim of optimal complexity. The cubic gives a certificate stronger than the preliminary $\varepsilon/(1 - 2\varepsilon)$ estimate.

II.29 Exact inversion and sharp recovery constants

Source 11, delivered Section 8.

The odd-target cubic factors as

$$\Psi_0(t) = 4t(1 - t)(1 - 3t/2).$$

The analytic inverse at zero therefore has a particularly explicit series.

Proposition II.29.1 (All-orders odd-target inverse). *For sufficiently small u ,*

$$D_0(u) = \sum_{n \geq 1} a_n u^n, \tag{II.29.1}$$

where

$$a_n = \frac{1}{n4^n} \sum_{j=0}^{n-1} \binom{n+j-1}{j} \binom{2n-j-2}{n-1-j} \left(\frac{3}{2}\right)^{n-1-j}. \tag{II.29.2}$$

In particular all the coefficients are positive, and the first eight are

$$\begin{array}{cccc} \frac{1}{4}, & \frac{5}{32}, & \frac{11}{64}, & \frac{475}{2048}, \\ \frac{1427}{4096}, & \frac{36645}{65536}, & \frac{123051}{131072}, & \frac{13661175}{8388608}. \end{array}$$

Proof. The inverse equation is

$$t = u \frac{1}{4(1-t)(1-3t/2)}.$$

The formal Lagrange coefficient identity gives

$$a_n = \frac{1}{n4^n} [z^{n-1}] (1-z)^{-n} (1-3z/2)^{-n}.$$

Expanding the two binomial series gives (II.29.2). Since $\Psi'_0(0) = 4$, the analytic inverse function theorem identifies the formal expansion with a convergent inverse in a neighborhood of zero. The exact rational substitution check through degree eight is included in the verification script. $\square$

No radius of convergence is needed for the quantitative theorem: D_0 is defined there by the unique root in $[0, 1/7]$. For numerical use, bisection on that interval gives a certified value without relying on a truncated asymptotic series. The beginning of the other branch is

$$D_1(u) = \frac{u}{4} + \frac{3u^2}{16} + \frac{u^3}{4} + O(u^4).$$

Proposition II.29.2 (Sharpness along vanishing-density families). *For either target type there are finite examples with $d(f) \rightarrow 0$ and*

$$d(f) = D_\kappa(\varepsilon(f)).$$

Consequently, no coefficient in the full small- u expansion can be decreased while keeping all preceding coefficients unchanged and an error of smaller order, uniformly over finite groups.

Proof. Take a one-point error in cyclic groups of increasing order, or a coset error of increasing index, using an involution when required. Let the error density be t . Once $t < 1/4$, the zero affine map is uniquely closest: any other affine map differs from it on at least half the domain and therefore differs from the error map on more than tN points. The exact defect is $\Psi_\kappa(t)$. For sufficiently large index this defect is below $1/9$, and equality in the inverse follows. Comparison along this sequence proves the coefficient assertion. $\square$

II.30 The largest non-affine energy in arbitrary large finite groups

Source 11, delivered Section 9.

We first prove the unrestricted-group part of Theorem F.16.3. The key additional ingredient is merely the discreteness of Hamming distance.

Proof of Theorem F.16.3 for $N \geq 34$. For $N \geq 34$,

$$\Psi_\kappa(1/N) \leq \Psi_0(1/34) = \frac{4290}{39304} < \frac{1}{9}.$$

Suppose a non-affine f had graph energy larger than the right-hand side of (F.16.7). Its defect would be smaller than $\Psi_\kappa(1/N)$. Theorem F.16.2 would imply $d(f) < 1/N$, impossible for a non-affine function on an N -point set.

At equality the same theorem gives $d(f) \leq 1/N$. Non-affinity forces $d(f) = 1/N$, and the equality description of the small-error theorem gives a one-point modification, of the required target order. Conversely, such a modification has exactly the asserted energy by Example II.24.3. It is non-affine because two distinct affine maps cannot differ on only one point when $N \geq 34$. $\square$

Corollary II.30.1 (Exact affine certification threshold). *For $N \geq 34$, if*

$$\mathcal{E}(f) > N^3 - 4N^2 + (10 + 2\kappa)N - (6 + 2\kappa),$$

then f is affine. The strict inequality cannot be weakened.

This assertion is a complete extremal statement for this finite problem, not just a stability estimate with an unspecified constant. It is also uniform in the structures and cardinalities of the target groups, except for the stated two-torsion distinction. The cutoff 34 is a convenient consequence of the correction radius, not a demonstrated optimal cutoff.

Remark W2 (The cutoff cannot simply be dropped; added at the third write, 6 October 2026). For general finite abelian groups the formula (F.16.7) can fail for $N < 34$ (it holds for some such N and fails for others); the smallest failure found has $N = 4$. Take $G = \mathbb{Z}/12\mathbb{Z}$, $W = \mathbb{Z}/2\mathbb{Z}$ (so $\kappa = 1$), and $f(x) = 1$ if $x \bmod 4 \in \{0, 1\}$, $f(x) = 0$ otherwise. Then f is not affine and

$$\mathcal{E}(f) = 1296 > 1288 = N^3 - 4N^2 + 12N - 8 \quad (N = 12).$$

Proof. Write $f = g \circ \pi$ with $\pi : \mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/4\mathbb{Z}$ the reduction and $g = \mathbf{1}_{\{0,1\}}$ on $\mathbb{Z}/4\mathbb{Z}$. Every additive quadruple of $\mathbb{Z}/12\mathbb{Z}$ maps to one of $\mathbb{Z}/4\mathbb{Z}$, and over each quadruple of $\mathbb{Z}/4\mathbb{Z}$ there are exactly $3^3 = 27$ (the lifts of x_1, x_2, x_3 are free and x_4 is then determined, and lies over the right point); whether f respects a quadruple depends only on its image. Hence $\mathcal{E}(f) = 27\mathcal{E}(g)$. By Lemma II.23.1, $\mathcal{E}(g) = \sum_h \sum_w r_h(w)^2$: the derivatives in directions 0 and 2 are constant ($\Delta_2 g \equiv 1$), contributing 16 each, and those in directions 1 and 3 take each value twice, contributing 8 each; so $\mathcal{E}(g) = 48$ and $\mathcal{E}(f) = 1296$. The derivative $\Delta_1 f$ takes both values ($\Delta_1 f(0) = 0$, $\Delta_1 f(1) = 1$), so f is not affine by Lemma II.23.2. *Exhaustive check at the write* (all maps, exact integers): the largest non-affine energies of maps $\mathbb{Z}/n\mathbb{Z} \rightarrow \mathbb{Z}/m\mathbb{Z}$ against the value of (F.16.7) are 48 against 40 for $\mathbb{Z}/4 \rightarrow \mathbb{Z}/2$, 85 against 77 for $\mathbb{Z}/5 \rightarrow \mathbb{Z}/2$, 231 against 223 for $\mathbb{Z}/7 \rightarrow \mathbb{Z}/2$, 162 against 126 for $\mathbb{Z}/6 \rightarrow \mathbb{Z}/3$, 513 against 489 for $\mathbb{Z}/9 \rightarrow \mathbb{Z}/3$ and 1296 against 1288 for $\mathbb{Z}/12 \rightarrow \mathbb{Z}/2$, while the formula holds for $\mathbb{Z}/10$, $\mathbb{Z}/14$, $\mathbb{Z}/15$ and $\mathbb{Z}/16 \rightarrow \mathbb{Z}/2$ (the intake found the same table, dossier `dossier116_GOWERS`). This does not contradict the theorem, which assumes $N \geq 34$, or its vector-space part ($\mathbb{Z}/12\mathbb{Z}$ and $\mathbb{Z}/4\mathbb{Z}$ are not vector spaces over a prime field). It shows that the cutoff cannot be removed for general groups and that its exact value is open (item II.21).

II.31 All finite vector-space sizes: a second proof

Source 11, delivered Section 10.

Let $V = \mathbb{F}_p^n$, $W = \mathbb{F}_p^m$, and $N = p^n$. Assume throughout this section that f is non-affine. The proof has no high-energy correction hypothesis. It explains exactly where prime-field structure helps with the small cardinalities.

II.31.1 Constant-derivative directions and quotienting

Define

$$H = \{h \in V : \Delta_h f \text{ is constant on } V\}.$$

By (II.23.2), H is a subspace, and the constant derivative values form a linear map $\lambda : H \rightarrow W$. Extend λ to a linear map $L : V \rightarrow W$ by extending a basis of H to a basis of V . The map $f - L$ is constant on each coset of H , so it factors through $\bar{f} : V/H \rightarrow W$.

Lemma II.31.1 (Energy under quotienting). *Let $q = |V/H|$. Then*

$$\frac{\mathcal{E}(f)}{N^3} = \frac{\mathcal{E}(\bar{f})}{q^3}. \quad (\text{II.31.1})$$

The map $\bar{f}$ has no constant derivative in a nonzero direction.

Proof. Every additive quadruple in the quotient lifts to exactly $|H|^3$ additive quadruples in V , and respect for the function is unchanged under this lift and subtraction of L . A nonzero constant-derivative direction in the quotient would lift to a direction in H , a contradiction. $\square$

II.31.2 The odd-characteristic histogram bound

Lemma II.31.2. *Let p be odd, q a positive power of p , and let a nonconstant q -term list in $\mathbb{F}_p^m$ have sum zero. Its value multiplicities satisfy*

$$\sum_w r(w)^2 \leq (q-2)^2 + 2. \quad (\text{II.31.2})$$

Equality holds precisely when the three nonzero multiplicities are $q-2, 1, 1$; the values are of the form $a, a+b, a-b$ with $b \neq 0$. For $q=3$ these three multiplicities are all one.

Proof. A multiplicity $q-1$ would give $(q-1)a+b=0$. Since $qa=0$, this forces $b=a$, contradicting nonconstancy. If a value has multiplicity $q-2$, the remaining two entries cannot be equal: equality to b would give $2(b-a)=0$. Thus they are distinct, and their sum must be $2a$. This gives the asserted equality configuration.

If the maximum multiplicity is at most $q-3$ and $q \geq 7$, convexity (or transferring units from smaller counts to larger counts) bounds the sum of squares by $(q-3)^2 + 3^2$, which is strictly less than $(q-2)^2 + 2$. For $q=5$, a cap of two gives at most $2^2 + 2^2 + 1^2 = 9 < 11$. For $q=3$, a nonconstant zero-sum list must have three distinct entries, which is exactly the asserted configuration. $\square$

For each derivative of $\bar{f}$, the sum of its values is zero, since translation permutes the domain. Applying the lemma in every nonzero direction and using the constant zero derivative in direction zero gives

$$\varepsilon(f) \geq \Psi_0(1/q). \quad (\text{II.31.3})$$

For real $q \geq 4$, the function $\Psi_0(1/q) = 4/q - 10/q^2 + 6/q^3$ is strictly decreasing; its derivative has numerator $-4q^2 + 20q - 18 < 0$ there. The only additional comparison needed among powers of an odd prime is $\Psi_0(1/3) > \Psi_0(1/9)$. Thus if $q < N$, the inequality in (II.31.3) is strictly stronger than $\varepsilon(f) \geq \Psi_0(1/N)$. This proves the odd-characteristic bound and shows that equality requires $H = \{0\}$ and equality in every nonzero-direction histogram.

II.31.3 Equality for odd characteristic

Assume first that $N \geq 7$. For each nonzero h , let $L(h)$ be the unique modal derivative value, attained $N-2$ times, and put $L(0) = 0$. For any h, k , the three modal identities in (II.23.2) exclude at most six choices of x . Since $N > 6$, there is a choice where all hold. Hence $L(h+k) = L(h) + L(k)$, and L is linear.

Put $g = f - L$. For every $h \neq 0$, exactly $N - 2$ values of x satisfy $g(x+h) = g(x)$. If $t_w = |g^{-1}(w)|$, summing these equality counts gives

$$\sum_w t_w^2 = N + (N - 1)(N - 2) = (N - 1)^2 + 1.$$

For a nonconstant map on N points, the largest sum of squared fiber sizes is $(N - 1)^2 + 1$, and equality forces the fiber sizes $N - 1, 1$. Indeed, among partitions of N into at least two positive parts, moving mass to the largest part increases the sum of squares until these two parts remain. Thus g is constant except at one point.

If $N = 5$, take a generator and write its derivative values as $a, a, a + b, a - b$, with $b \neq 0$. After subtracting the linear map $x \mapsto ax$, the derivative is zero except for one b and one $-b$. Going around the five-cycle, the function has two values on complementary cyclic intervals. If either interval is a singleton, the required conclusion follows. Otherwise their sizes are two and three; the step-two derivative then has multiplicities 1, 2, 2, whose squared sum is $9 < 11$, contradicting equality.

If $N = 3$, subtract the affine map interpolating any two of the three points. A non-affine map then differs at exactly the remaining point. This completes the equality proof for every odd characteristic.

II.31.4 Characteristic two

If $p = 2$, each derivative obeys $\Delta_h f(x + h) = \Delta_h f(x)$, so for $h \neq 0$ every histogram multiplicity is even. On a quotient of size q , a nonconstant histogram therefore has squared sum at most $(q - 2)^2 + 4$, attained by multiplicities $q - 2, 2$. A non-affine map cannot have quotient size two, since every map from $\mathbb{F}_2$ to W is affine. Thus $q \geq 4$, and

$$\varepsilon(f) \geq \Psi_1(1/q).$$

The function $\Psi_1(1/q)$ is decreasing for real $q \geq 8$, and $\Psi_1(1/4) > \Psi_1(1/8)$. Hence it is strictly decreasing along the powers of two that arise here. Equality at the N -point bound again forces $H = \{0\}$ and extremal histograms in every nonzero direction.

For $N \geq 8$, the modal values have exactly two exceptions, so the same six-exception argument makes them a linear map. The fiber-square argument then gives a one-point modification. For $N = 4$, interpolate a linear-affine map through $0, e_1, e_2$ and subtract it. Only the point $e_1 + e_2$ remains, so every non-affine map has the required form. This proves the remaining part of Theorem F.16.3.

II.32 Partial domains and the interface with Proposition 6.1

Source 11, delivered Section 11.

Gowers's frequency-selection map is naturally defined on a subset rather than the entire group. The full-energy normalization, not a normalization by the subset size, gives a clean extension.

Theorem II.32.1 (Canonical affine recovery from a partial graph). *Let $B \subseteq G$, $\phi : B \rightarrow W$, and let Γ_B be its graph. Suppose*

$$E(\Gamma_B) \geq (1 - \varepsilon)N^3, \quad 0 \leq \varepsilon < 1/9.$$

There is a unique affine map $\ell : G \rightarrow W$ satisfying

$$|\{x \in B : \phi(x) \neq \ell(x)\}| \leq D_\kappa(\varepsilon)N. \tag{II.32.1}$$

Moreover,

$$|G \setminus B|/N \leq 1 - (1 - \varepsilon)^{1/3}. \quad (\text{II.32.2})$$

The same affine map is recovered by applying the full-domain theorem to any completion of ϕ .

Proof. Set $\beta = |B|/N$. A graph with $|B|$ points has energy at most $|B|^3$, so $\beta^3 \geq 1 - \varepsilon$, proving (II.32.2). Extend ϕ arbitrarily to $f : G \rightarrow W$. Every respected quadruple of the partial graph is also a respected quadruple of the full graph, so $\varepsilon(f) \leq \varepsilon$. Theorem F.16.2 and monotonicity of D_κ give an affine map satisfying (II.32.1).

If two affine maps satisfy this inequality, their agreement set has size at least $(\beta - 2D_\kappa(\varepsilon))N$. Here $D_\kappa(\varepsilon) < 1/7$, while $\beta > (8/9)$, unless $\varepsilon = 0$, when the conclusion is immediate. Thus

$$\beta - 2D_\kappa(\varepsilon) > \frac{8}{9} - \frac{2}{7} = \frac{38}{63} > \frac{1}{2}.$$

Distinct affine maps cannot agree on more than half the domain. This proves uniqueness and independence from the completion. $\square$

The error in (II.32.1) is normalized by N , not by $|B|$. The statement requires a graph energy close to N^3 . It does not assert that an arbitrary sparse partial homomorphism with energy close to $|B|^3$ extends to an affine map on the ambient group.

Corollary II.32.2 (A high-correlation refinement of the Section 6 output). *Let N be an odd prime, $F : \mathbb{Z}/N\mathbb{Z} \rightarrow \mathbb{C}$ satisfy $|F| \leq 1$, and let $\phi : B \rightarrow \mathbb{Z}/N\mathbb{Z}$ for $B \subseteq \mathbb{Z}/N\mathbb{Z}$. Use the unnormalized Fourier transform $\hat{u}(r) = \sum_x u(x) \exp(-2\pi i r x/N)$ and put $\partial_h F(x) = F(x)\overline{F(x-h)}$. Suppose $\alpha > 0$ and*

$$\sum_{h \in B} |\widehat{\partial_h F}(\phi(h))|^2 \geq \alpha N^3, \quad \alpha^4 > 8/9.$$

Then the frequency-selection map ϕ agrees on B with a unique affine map, except on at most

$$D_0(1 - \alpha^4)N$$

points. Also $|B| \geq \alpha^{4/3}N$.

Proof. Proposition 6.1 gives at least $\alpha^4 N^3$ respected quadruples, and the graph identity in Lemma 6.2 converts this to the hypothesis of Theorem II.32.1. The target $\mathbb{Z}/N\mathbb{Z}$ has no nonzero involution. $\square$

[write] 6 October 2026; sources 28–33. Source 31's Theorem II.108.1 treats selected graphs on every finite abelian group, partial graphs included and with no amplitude bound: graph errors at most ε and quadratic infidelity at most $(1 - \sqrt{1 - 2\varepsilon})/2$ for score $1 - \varepsilon$, $\varepsilon \leq 1/400$, both separately sharp.

[write] 7 October 2026; sources 54–59. Source 59 lowers the condition $\alpha^4 > 8/9$ to $\alpha^4 > 7/9$ for targets without involutions (its Corollary II.251.4 and the paragraph after it).

The approximation error begins

$$\left(\frac{1 - \alpha^4}{4} + \frac{5}{32}(1 - \alpha^4)^2 + \frac{11}{64}(1 - \alpha^4)^3 + O((1 - \alpha^4)^4) \right) N.$$

This corollary genuinely strengthens the structural information obtained from that proposition in its high-correlation range. It does not improve the exponent four in the original energy lower bound.

In the usual small- α inverse-theorem regime, its hypothesis is unavailable; this is the precise reason it does not automatically improve the global Szemerédi bound.

[write] **6 October 2026; sources 11–14.** *Scope (intake observation, checked at the write).* The hypothesis that N is an odd prime is not needed for the argument. The formal `proposition_6_1` (`Sections06_07.lean`, line 37; exact companion `proposition_6_1_holds`, `Proofs05_10.lean`, line 505, at HEAD `e1db00e99`) assumes only $N \geq 1$, $0 < \alpha$, $|F| \leq 1$ and the displayed Fourier hypothesis with the unnormalized transform and the same derivative $F(x)F(x-h)$, and concludes at least $\alpha^4 N^3$ respected quadruples; the graph identity of Lemma 6.2 and Theorem II.32.1 hold for every finite abelian group. So the corollary holds for every odd N as stated, and for even N with D_1 in place of D_0 (then $\mathbb{Z}/N\mathbb{Z}$ contains the involution $N/2$). Not formalized. *Relation to source 02.* Source 02’s Theorem II.2.1 strengthens the energy *lower bound* of Proposition 6.1; this corollary uses that bound’s output in the range $\alpha^4 > 8/9$. The two results are complementary and share no theorem.

II.33 Source 12’s chapter: Freiman extension and Bohr radii in Section 10 [write]

Added in the third write, 6 October 2026.

The next nine sections are source 12, a manuscript on the second half of Gowers’s Section 10, Lemmas 10.10–10.14: extending a Freiman homomorphism from a dense subset of a Bohr set to a smaller Bohr set, choosing the Bohr radii, and the local model of Theorem 10.13. *Logically the chapter continues source 07’s chapter* (Section II.11), which treats Lemma 10.9, the step just before: together they cover Section 10 from Lemma 10.6 to Corollary 10.14, and they share no statement. It is printed here, at the end of Part II, because no number printed earlier may move. Its notation follows the Section 10 table at the head of this Part. Its introduction is Section F.17; its formalization interfaces, checks, questions and parameter ledger are in Part V (Sections V.50–V.52). Its Appendix A, the fixed-domain repair of the extraction input, is printed at the end of this chapter, beside the theorem that uses it.

What is proved, in one paragraph. If $S \subseteq B$ misses at most a fraction ε of B and every shift in C loses at most a fraction κ of B , a Freiman homomorphism on S induces one on C whenever $3\varepsilon + 2\kappa < 1$, and every rational boundary point has a counterexample (Theorems II.35.2, II.35.3); the collision-to-mode bound $1 - \chi(\eta)$ is sharp (Lemma II.36.1); a finite pigeonhole among m radii selects an ambient radius with overlap $4^{-d/m}$ (Theorem II.37.2); a fourth-moment estimate enlarges the initial radius from $2^{-148}\alpha^{18}/\pi$ to $2^{-74}\alpha^9/\pi$ with the same spectrum (Theorem II.38.1); and Theorem 10.13 holds on a Bohr set of radius at least $\alpha^{19}/(3\pi 2^{149})$ (Theorem II.40.2, Corollary II.40.3). The component-extraction input (Lemmas 10.3–10.9) is borrowed, with its hypotheses stated (Lemma II.40.1); sharpness is universal over groups, not for prime cyclic groups with the same target (the source’s non-claims).

Relation to the formal statements (checked at HEAD `e1db00e99`). `lemma_10_10` (`Section10.lean`, line 362; `lemma_10_10_holds`, `Proofs10Bohr.lean`, line 277), `corollary_10_11` (line 374; `Proofs10Bohr.lean`, line 607), `lemma_10_12` (line 401; `Proofs10Induced.lean`, line 145), `theorem_10_13` (line 473; `Proofs10Main.lean`, line 1292) and `corollary_10_14` (line 501; `Proofs10ProgressionLinearity.lean`, line 133) are exact companions; `theorem_10_13` assumes $\alpha \leq 1/6$ and exports $\alpha^6|X|/20000$, while its proof already establishes the fraction $5\alpha^5/98304$ (note in Section F.17). Nothing in this chapter is formalized.

II.34 Translation loss and intersections along a tree

Source 12, delivered Section 2.

Throughout the algebraic parts, G and H are abelian groups written additively. The ambient group G need not be finite, but $B \subseteq G$ is finite and nonempty. For $t \in G$, $B - t = \{x : x + t \in B\}$. All missing densities are measured relative to $|B|$, not $|G|$.

Definition II.34.1. A function $\phi : S \rightarrow H$ is a *Freiman homomorphism of order two* if

$$a + b = c + d, \quad a, b, c, d \in S \quad \implies \quad \phi(a) + \phi(b) = \phi(c) + \phi(d).$$

Repeated elements are permitted. A Freiman homomorphism of order r preserves all equalities between sums of r elements of its domain.

Lemma II.34.2 (Elementary properties of translation loss). *For every $u, v \in G$,*

$$b_B(0) = 0, \quad b_B(-u) = b_B(u), \quad 0 \leq b_B(u) \leq 1, \quad b_B(u + v) \leq b_B(u) + b_B(v).$$

Also $|B \triangle (B - u)| = 2b_B(u)|B|$.

Proof. Finite translates have equal cardinality, so each one-sided difference has half the cardinality of the symmetric difference. Translation gives symmetry. For subadditivity, use

$$B \setminus (B - u - v) \subseteq [B \setminus (B - u)] \cup [(B - u) \setminus (B - u - v)]$$

and count. The other claims follow directly from the definition. $\square$

Lemma II.34.3 (Intersection bound along a tree). *Let $t_0, \dots, t_m \in G$, and let T be a tree on the labels $0, \dots, m$. Suppose $S_i \subseteq B$ and $|B \setminus S_i| \leq \varepsilon_i |B|$. Then*

$$\left| \bigcap_{i=0}^m (S_i - t_i) \right| \geq \left(1 - \sum_{i=0}^m \varepsilon_i - \sum_{\{i,j\} \in E(T)} b_B(t_i - t_j) \right) |B|. \quad (\text{II.34.1})$$

The right side may be negative; its useful content is when it is positive.

Proof. Root the tree at 0 and add its vertices in an order in which each parent precedes its children. Before adding a child j with parent i , the current intersection of ambient translates is contained in $B - t_i$. Imposing membership in $B - t_j$ therefore removes at most

$$|(B - t_i) \setminus (B - t_j)| = b_B(t_j - t_i)|B|$$

points. Starting with a translate of B proves the bound with all $S_i = B$. Replacing $B - t_i$ by $S_i - t_i$ removes at most $\varepsilon_i |B|$ further points. Summing these losses proves (II.34.1). $\square$

[write] **6 October 2026; sources 28–33.** Source 30's Lemma II.131.2 is this lemma (with losses counted in points), proved again without citation.

This formulation records the geometry of the configuration. One pays for the edges of a spanning tree, and can choose the tree of least total translation loss. It avoids charging every displacement independently from a single base point.

Corollary II.34.4 (Two and three translated copies). *If $S \subseteq B$ and $|S| \geq (1 - \varepsilon)|B|$, then*

$$|S \cap (S - t)| \geq (1 - 2\varepsilon - b_B(t))|B|, \quad (\text{II.34.2})$$

$$|S \cap (S - c) \cap (S - c - d)| \geq (1 - 3\varepsilon - b_B(c) - b_B(d))|B|. \quad (\text{II.34.3})$$

In the second inequality, $b_B(c) + b_B(d)$ can be replaced by

$$\tau_B(c, d) = \min\{b_B(c) + b_B(d), b_B(c) + b_B(c + d), b_B(d) + b_B(c + d)\}.$$

Proof. Apply lemma II.34.3 to the vertices $0, t$ or $0, c, c + d$. The three possible trees on the latter vertices give the stated minimum. $\square$

II.35 A sharp theorem for induced Freiman homomorphisms

Source 12, delivered Section 3.

Lemma II.35.1 (The canonical difference map). *If $S \neq \emptyset$ and $\phi : S \rightarrow H$ is a Freiman homomorphism, there is a unique map $D : S - S \rightarrow H$ satisfying*

$$D(y - x) = \phi(y) - \phi(x) \quad (x, y \in S).$$

It obeys $D(0) = 0$ and $D(-t) = -D(t)$. If $x, x + c, x + c + d \in S$, then

$$D(c) + D(d) = D(c + d). \quad (\text{II.35.1})$$

Proof. If $y - x = v - u$, then $y + u = v + x$, so the Freiman property gives $\phi(y) - \phi(x) = \phi(v) - \phi(u)$. This proves well-definedness and uniqueness. Swapping the ordered pair gives oddness. The three-point identity follows by telescoping the two successive increments. $\square$

Theorem II.35.2 (Sharp density–overlap extension). *Let $S \subseteq B$ with $|S| \geq (1 - \varepsilon)|B|$, and let $C \subseteq G$ be symmetric with $0 \in C$. Suppose*

$$b_B(c) \leq \kappa \quad (c \in C), \quad 3\varepsilon + 2\kappa < 1.$$

Then, for every Freiman homomorphism $\phi : S \rightarrow H$:

- (i) $C + C \subseteq S - S$;
- (ii) *its canonical difference map satisfies $D(c) + D(d) = D(c + d)$ for every $c, d \in C$;*
- (iii) $D|_C$ *is a Freiman homomorphism of order two and $D(0) = 0$.*

More precisely, every ordered pair $c, d \in C$ has at least $(1 - 3\varepsilon - 2\kappa)|B|$ witnesses x with $x, x + c, x + c + d \in S$.

Proof. The witness count is (II.34.3). It is positive, so for each c, d a witness exists. Its endpoints show that $c + d \in S - S$, and lemma II.35.1 gives (ii). If $c_1 + c_2 = c_3 + c_4$, with all $c_i \in C$, both sums of D -values equal the same well-defined value $D(c_1 + c_2)$. This proves (iii). $\square$

[write] **6 October 2026; sources 28–33.** Source 30's Theorem II.131.3 needs $4\varepsilon + 3\kappa < 1$ at order two, which implies the condition here; with source 30's fixed-radius estimate this theorem gives the larger fixed radius of the note after its Corollary II.132.2.

Only losses for shifts in C are required. Imposing the same loss bound on $C+C$ would unnecessarily shrink the eventual Bohr radius. The proof also explains why one does not need to find one base point supporting two entire decompositions of a sum: the canonical difference at their common endpoint makes two separate triangles sufficient.

Theorem II.35.3 (Sharpness on the whole boundary). *Every rational pair (ε, κ) with $0 \leq \varepsilon \leq 1/3$, $0 \leq \kappa \leq 1/2$, and $3\varepsilon + 2\kappa = 1$ occurs in a finite example satisfying the size and overlap hypotheses of theorem II.35.2 but not its conclusion.*

Proof. Let K be a finite abelian group of size m , choose $T \subseteq K$ of size t , $0 \leq t \leq m$, and put

$$G = \mathbb{Z}/3\mathbb{Z} \times K, \quad S = \{0, 1\} \times K, \quad B = S \cup (\{2\} \times T), \quad C = \mathbb{Z}/3\mathbb{Z} \times \{0\}.$$

Define $\phi : S \rightarrow \mathbb{Z}$ by $\phi(a, k) = a$ using the representatives $a = 0, 1$. This is a Freiman homomorphism: the sum of two first coordinates is one of $0, 1, 2$, and equality modulo 3 of such sums is equality as integers.

We have $|B| = 2m + t$, $|S| = 2m$, and, for either nonzero $c \in C$, $|B \cap (B - c)| = m + 2t$. Therefore

$$\varepsilon = \frac{t}{2m + t}, \quad \kappa = \frac{m - t}{2m + t}, \quad 3\varepsilon + 2\kappa = 1.$$

Writing $e = (1, 0)$, the induced values are $D(e) = 1$, $D(2e) = -1$, and $D(0) = 0$. But $e + e = 2e + 0$, whereas $1 + 1 \neq -1 + 0$. Thus $D|_C$ is not a Freiman homomorphism. Finally, every rational $\varepsilon \in [0, 1/3]$ is obtained by taking $t/m = 2\varepsilon/(1 - \varepsilon)$ and clearing denominators. $\square$

Remark II.35.4 (Scope of sharpness). This is a sharp universal statement over abelian ambient and target groups. It does not assert an optimal threshold when both groups are the same prime cyclic group. The example uses a target in which three times a nonzero value need not vanish. Symmetric ambient examples are also available by using the two slices $\{-1, 1\} \times K$ and a symmetric subset of the zero slice, with the corresponding integer-valued map.

Corollary II.35.5 (Global extension above two thirds). *If G is finite, $|S| > 2|G|/3$, and $\phi : S \rightarrow H$ is a Freiman homomorphism, then there are a unique homomorphism $D : G \rightarrow H$ and a unique $h_0 \in H$ such that*

$$\phi(x) = h_0 + D(x) \quad (x \in S).$$

The constant $2/3$ cannot be reduced uniformly over G, H .

Proof. Take $B = C = G$ and $\kappa = 0$. The canonical difference map is additive on all of G . Fix $x_0 \in S$ and put $h_0 = \phi(x_0) - D(x_0)$. The difference identity gives the formula for all $x \in S$. Since $S - S = G$, any affine extension has this same difference map and then the same constant. Set $t = m$ in theorem II.35.3 for the endpoint obstruction. $\square$

Theorem II.35.6 (Higher-order local extension). *Let $r \geq 2$ be an integer. Under the same set hypotheses, replace the numerical condition by $3\varepsilon + r\kappa < 1$. Then $rC \subseteq S - S$, and the canonical difference map satisfies*

$$D(c_1 + \cdots + c_j) = D(c_1) + \cdots + D(c_j) \quad (1 \leq j \leq r, \ c_i \in C).$$

In particular, $D|_C$ is a Freiman homomorphism of order r .




Figure II.35.1: The universal sufficient region is below the line $3\varepsilon + 2\kappa = 1$. The construction in theorem II.35.3 supplies counterexamples densely along the boundary. The diagram records an exact theorem, not an empirical fit.

Proof. Subadditivity gives $b_B(t) \leq r\kappa$ for $t \in rC$. The pair bound is positive since $1 - 2\varepsilon - r\kappa > 0$, so $rC \subseteq S - S$. For a partial sum $u = c_1 + \dots + c_{j-1}$, the triangle loss is at most $b_B(u) + b_B(c_j) \leq j\kappa$. Thus (II.35.1) applies and proves the identity inductively. Two equal sums of r elements have the same canonical difference value. We make no optimality claim for this higher-order numerical condition. $\square$

II.36 Correction from concentrated directional increments

Source 12, delivered Section 4.

The exact Freiman hypothesis can be weakened in a way that separates algebraic errors from missing points. This is useful for local repair, although a global average over all additive quadruples does not by itself give the uniform directional hypothesis below.

For any function $\phi : S \rightarrow H$ and any $t \in S - S$, let

$$E_t = \{x \in S : x + t \in S\}, \quad p_t(h) = \frac{|\{x \in E_t : \phi(x+t) - \phi(x) = h\}|}{|E_t|}.$$

These distributions have finite support even when H is infinite. Their collision probabilities are

$$Q_t = \sum_{h \in H} p_t(h)^2.$$

Equivalently, Q_t is the fraction of ordered pairs of t -edges with the same increment label.

Lemma II.36.1 (Sharp collision-to-majority bound). *For $0 \leq \eta < 1/2$ define*

$$\chi(\eta) = \frac{1 - \sqrt{1 - 2\eta}}{2}.$$

If a finite probability distribution p has $\sum p_i^2 \geq 1 - \eta$, its largest atom is unique and has mass at least $1 - \chi(\eta)$. This bound is sharp.

Proof. Let $a = \max p_i$. Since $\sum p_i^2 \leq a$, we have $a \geq 1 - \eta > 1/2$. Also

$$\sum p_i^2 \leq a^2 + (1 - a)^2.$$

Solving $a^2 + (1 - a)^2 \geq 1 - \eta$ on the interval $a > 1/2$ gives the asserted lower bound. A two-atom distribution with masses $1 - \chi(\eta)$ and $\chi(\eta)$ attains equality. $\square$

[write] **6 October 2026; sources 11–14.** *Credit.* This is the sharpest of the report's collision-to-mode bounds: source 07's Proposition II.19.1 and source 11's Lemma II.28.1 (which uses $\sum p_i^2 \leq \max p_i$, mode at least $1 - \eta$) are weaker, since $\chi(\eta) \leq \eta$. Source 14's envelope $\tau(\varepsilon) = (1 - \sqrt{1 - 4\varepsilon})/2$ is the same function, $\tau(\varepsilon) = \chi(2\varepsilon)$, in a different problem.

Theorem II.36.2 (Directional repair). *Let $S \subseteq B$, $C = -C$, $0 \in C$, $|S| \geq (1 - \varepsilon)|B|$, and $b_B(c) \leq \kappa$ for $c \in C$. Suppose*

$$3\varepsilon + 2\kappa + 3\chi(\eta) < 1, \quad 0 \leq \eta < 1/2. \quad (\text{II.36.1})$$

Assume $Q_t \geq 1 - \eta$ for every $t \in C + C$. Then the majority increment $D(t)$ is well defined on $C + C$, and

$$D(c) + D(d) = D(c + d) \quad (c, d \in C).$$

Consequently $D|_C$ is an exact Freiman homomorphism, $D(0) = 0$, and at least $1 - \chi(\eta)$ of the t -edges carry label $D(t)$ for every $t \in C + C$.

Proof. The strict budget first implies $3\varepsilon + 2\kappa < 1$, so the geometric triangle bound shows $C + C \subseteq S - S$ without any algebraic assumption on ϕ . Thus the distributions in the hypothesis exist. By lemma II.36.1, each majority label is unique and at most $\chi(\eta)|E_t|$ of its edges are bad.

Fix $c, d \in C$ and put $s = c + d$. There are at least $(1 - 3\varepsilon - 2\kappa)|B|$ points x with $x, x + c, x + s \in S$. Among these witnesses, at most $\chi(\eta)|E_c|$ have a bad first edge, at most $\chi(\eta)|E_d|$ have a bad second edge (translate its starting point by c), and at most $\chi(\eta)|E_s|$ have a bad direct edge. Each $|E_t| \leq |B|$, so (II.36.1) leaves a witness with all three edges good. Telescoping its actual ϕ -increments proves the additive identity. The remaining conclusions follow as in theorem II.35.2; the zero increment is identically zero. $\square$

Remark II.36.3 (The exact nonuniform budget). For a fixed triangle, one may instead require

$$1 - 3\varepsilon - \tau_B(c, d) > \frac{\chi(\eta_c)|E_c| + \chi(\eta_d)|E_d| + \chi(\eta_{c+d})|E_{c+d}|}{|B|},$$

where $Q_t \geq 1 - \eta_t$ and each $\eta_t < 1/2$. This follows from precisely the same witness count. The three equal error allowances in (II.36.1) are a convenient uniform specialization, not an asserted optimal repair threshold.

Corollary II.36.4 (Global affine repair). *Take $B = C = G$ finite, $|S| \geq (1 - \varepsilon)|G|$, and assume $3\varepsilon + 3\chi(\eta) < 1$ and $Q_t \geq 1 - \eta$ for every $t \in G$. Then the majority increments give a homomorphism $D : G \rightarrow H$, and some $h_0 \in H$ satisfies*

$$|\{x \in S : \phi(x) = h_0 + D(x)\}| \geq [1 - \chi(\chi(\eta))]|S|.$$

In particular the weaker, simpler proportion $1 - \chi(\eta)$ also holds.

Proof. The repair theorem makes D additive everywhere. Put $F(x) = \phi(x) - D(x)$ on S . Every ordered pair $(x, y) \in S^2$ appears in exactly one direction $t = y - x$. It has $F(y) \neq F(x)$ exactly when its directional edge is bad. Hence

$$|\{(x, y) \in S^2 : F(x) \neq F(y)\}| \leq \chi(\eta) \sum_t |E_t| = \chi(\eta) |S|^2.$$

The distribution of F therefore has collision probability at least $1 - \chi(\eta)$. A second application of lemma II.36.1 gives the result. $\square$

The error parameter is conditional on a direction. Replacing this by the assertion that a $1 - \eta$ fraction of all additive quadruples is respected is a different problem: directions with few representations can have large conditional error. The component-extraction stage of Gowers's proof is designed to cope with this difficulty. Our repair theorem does not silently remove that stage.

II.37 A finite selection argument for Bohr radii

Source 12, delivered Section 5.

Let G now be a finite abelian group of order N , and identify characters with homomorphisms $\gamma : G \rightarrow \mathbb{T} = \mathbb{R}/\mathbb{Z}$. Write $\|u\|_{\mathbb{T}}$ for distance to the nearest integer. For a finite set Γ of d characters and $r > 0$, set

$$\mathcal{B}(\Gamma, r) = \{x \in G : \|\gamma(x)\|_{\mathbb{T}} \leq r \text{ for every } \gamma \in \Gamma\}.$$

These are closed Bohr sets, matching the weak inequalities in the repository. The case $\Gamma = \emptyset$ gives G and reduces to the global extension theorem; below $d \geq 1$.

Lemma II.37.1 (Elementary size and doubling bounds). *For $0 < r \leq 1/2$,*

$$|\mathcal{B}(\Gamma, r)| \geq \frac{N}{\lceil 1/r \rceil^d} \geq (r/2)^d N, \quad (\text{II.37.1})$$

$$|\mathcal{B}(\Gamma, r)| \leq 4^d |\mathcal{B}(\Gamma, r/2)|. \quad (\text{II.37.2})$$

Also $\mathcal{B}(\Gamma, r) + \mathcal{B}(\Gamma, s) \subseteq \mathcal{B}(\Gamma, r + s)$.

Proof. Map x to its d character values in $\mathbb{T}^d$. Divide each circle into $\lceil 1/r \rceil$ half-open intervals of equal length at most r . Some product cell has at least $N/\lceil 1/r \rceil^d$ preimages. Subtract a fixed preimage from all others in that cell; the differences belong to $\mathcal{B}(\Gamma, r)$. The weaker bound follows from $\lceil 1/r \rceil \leq 2/r$.

For (II.37.2), represent all character values of $\mathcal{B}(\Gamma, r)$ in $[-r, r]$. Partition this cube into 4^d cells of side length $r/2$, assigning boundary points consistently. In each occupied cell choose one preimage. Subtraction from that preimage injects the cell's preimages into $\mathcal{B}(\Gamma, r/2)$. Sum over the cells. The final assertion is the triangle inequality on $\mathbb{T}$. $\square$

Theorem II.37.2 (One-scale radius selection). *Let $0 < \rho \leq 1/2$ and let $m \geq 1$ be an integer. Put $s = \rho/(2m)$. There is*

$$r \in \left\{ \frac{\rho}{2} + js : 1 \leq j \leq m \right\} \subseteq (\rho/2, \rho]$$

such that, for $B = \mathcal{B}(\Gamma, r)$,

$$|B \cap (B - c)| \geq 4^{-d/m} |B| \quad (c \in \mathcal{B}(\Gamma, s)). \quad (\text{II.37.3})$$

More generally, replace 4^d by the actual ratio

$$R = \frac{|\mathcal{B}(\Gamma, \rho)|}{|\mathcal{B}(\Gamma, \rho/2)|}$$

and replace $4^{-d/m}$ by $R^{-1/m}$.

Proof. Set $r_j = \rho/2 + js$, $0 \leq j \leq m$, and $a_j = |\mathcal{B}(\Gamma, r_j)|$. The product of a_{j-1}/a_j is $a_0/a_m = R^{-1} \geq 4^{-d}$. Therefore some ratio is at least $R^{-1/m} \geq 4^{-d/m}$. Choose that j and $r = r_j$. If $x \in \mathcal{B}(\Gamma, r - s)$ and $c \in \mathcal{B}(\Gamma, s)$, both x and $x + c$ belong to B . Thus the inner set is contained in $B \cap (B - c)$, proving the assertion. $\square$

This proves the needed overlap at one scale with a finite pigeonhole argument. It does not assert that B is regular at every smaller relative change of radius in the standard sense. Full regularity is available from the classical theory [4], but is unnecessary here.

Corollary II.37.3 (A replacement for Corollary 10.11). *Given Γ , $d \geq 1$, and $0 < \rho \leq 1/2$, there exists $r \in (\rho/2, \rho]$ such that every $S \subseteq B = \mathcal{B}(\Gamma, r)$ with $|S| \geq 7|B|/8$ and every Freiman homomorphism $\phi : S \rightarrow H$ induce a Freiman homomorphism on*

$$C = \mathcal{B}(\Gamma, \rho/(8d)).$$

In fact $C + C \subseteq S - S$ and the induced difference map is additive on pairs from C .

Proof. Take $m = 4d$ in theorem II.37.2. Then $\kappa = 1 - 1/\sqrt{2}$ is admissible, and

$$3 \cdot \frac{1}{8} + 2 \left(1 - \frac{1}{\sqrt{2}}\right) = \frac{19}{8} - \sqrt{2} < 1.$$

Apply theorem II.35.2. $\square$

[write] **6 October 2026; sources 11–14.** *Scope.* As the remark “Why the order of selection matters” in Section II.37 and the caption of Table F.17.1 say, the ambient radius r is *selected* in $(\rho/2, \rho]$ before the dense subset is chosen. So this corollary is an alternative to Corollary 10.11 with a changed quantifier order, not a replacement for it at a prescribed radius; the abstract’s word “replaces” (Section F.17) should be read this way. *Observation (intake, checked at the write):* the hypotheses $C = -C$ and $0 \in C$ of Theorem II.35.2 are not used by its proof: for $c, d \in C$ the witness count gives $x, x + c, x + c + d \in S$, hence $c, d, c + d \in S - S$, and the conclusions follow; $D(0) = 0$ holds since $S \neq \emptyset$.

[write] **6 October 2026; sources 28–33.** At a prescribed radius $\rho \leq 1/4$ source 30 gives $\rho/(8(3^r - 1))$ with the same margins (Corollary II.132.1; r here is d), and shows that a fixed-radius estimate must depend exponentially on the rank (Corollary II.135.5); this corollary’s $\rho/(8d)$ uses the selected radius.

Corollary II.37.4 (General density and extension order). *Let $0 \leq \varepsilon < 1/3$ and $q \geq 2$. If the integer m satisfies*

$$3\varepsilon + q(1 - 4^{-d/m}) < 1,$$

the selected radius in theorem II.37.2 makes every Freiman homomorphism on a $(1 - \varepsilon)$ -dense subset of B induce a Freiman homomorphism of order q on $\mathcal{B}(\Gamma, \rho/(2m))$. In particular, one can take any integer

$$m > \frac{d \log 4}{-\log(1 - (1 - 3\varepsilon)/q)}.$$

Proof. Use theorem II.35.6 with $\kappa = 1 - 4^{-d/m}$. Rearranging the strict inequality gives the displayed sufficient choice. $\square$

Remark II.37.5 (Why the order of selection matters). The radius r depends only on Γ, ρ, m , not on the later subset or map. This is exactly the quantifier order required in the application. If a dense subset of $\mathcal{B}(\Gamma, \rho)$ has already been fixed, shrinking the ambient radius can concentrate all its missing points in the smaller set. Its relative density need not survive. We therefore select r before invoking component extraction.

Example II.37.6 (Why a factor of the rank is natural). In $G = (\mathbb{Z}/L\mathbb{Z})^d$, take the d coordinate characters and $\rho < 1/4$. For large L , the Bohr set is a product of intervals of length about $2\rho L$. A shift whose every coordinate is about sL has overlap ratio tending to

$$\left(1 - \frac{s}{2\rho}\right)^d.$$

Keeping this uniformly above a fixed positive constant requires $s = O(\rho/d)$. Thus the inverse-rank scale in the overlap method is necessary for this family; the example does not assert optimality of the constants 6 or 8.

II.38 Fourier smoothing with the full fourth moment

Source 12, delivered Section 6.

We use unnormalized Fourier transforms and convolutions in this section:

$$\widehat{g}(\gamma) = \sum_{x \in G} g(x) e^{-2\pi i \gamma(x)}, \quad (f * g)(x) = \sum_{y \in G} f(y) g(x - y).$$

Thus $g(x) = N^{-1} \sum_{\gamma} \widehat{g}(\gamma) e^{2\pi i \gamma(x)}$ and $\sum_{\gamma} |\widehat{g}(\gamma)|^2 = N \sum_x |g(x)|^2$.

Theorem II.38.1 (Fourth-moment smoothing). *Let $0 \leq g(x) \leq M$ and $\sum_x g(x) = \alpha MN$, where $M > 0$ and $0 < \alpha \leq 1$. Put*

$$L = \alpha^3 M^4 N^3, \quad h = g * \widetilde{g} * g * \widetilde{g}, \quad \widetilde{g}(x) = g(-x),$$

and, for $\lambda > 0$, let

$$\Gamma = \{\gamma : |\widehat{g}(\gamma)| \geq \lambda MN\}, \quad d = |\Gamma|.$$

Then, for every $\rho > 0$,

$$d \leq \frac{\alpha}{\lambda^2}, \tag{II.38.1}$$

$$\sum_{\gamma} |\widehat{g}(\gamma)|^4 \leq LN, \tag{II.38.2}$$

$$\sup_x |h(x+t) - h(x)| \leq \left(2\pi\rho + \frac{2\lambda^2}{\alpha^2}\right) L \quad (t \in \mathcal{B}(\Gamma, \rho)). \tag{II.38.3}$$

More precisely, if $\xi = (LN)^{-1} \sum_{\gamma} |\widehat{g}(\gamma)|^4$, then $\alpha \leq \xi \leq 1$ and $2\pi\rho$ in (II.38.3) can be replaced by $2\pi\rho\xi$.

Proof. Parseval and $g^2 \leq Mg$ give

$$\sum_{\gamma} |\widehat{g}(\gamma)|^2 = N \sum_x g(x)^2 \leq \alpha M^2 N^2.$$

Each element of Γ contributes at least $\lambda^2 M^2 N^2$, proving (II.38.1). Since $|\widehat{g}(\gamma)| \leq \alpha MN$,

$$\sum_{\gamma} |\widehat{g}(\gamma)|^4 \leq (\alpha MN)^2 \sum_{\gamma} |\widehat{g}(\gamma)|^2 \leq \alpha^3 M^4 N^4 = LN.$$

The zero-frequency term alone is $\alpha^4 M^4 N^4 = \alpha LN$, proving the lower bound for ξ .

The convolution identity is $\widehat{h}(\gamma) = |\widehat{g}(\gamma)|^4$. On Γ , membership of t in the Bohr set gives

$$|e^{2\pi i \gamma(t)} - 1| \leq 2\pi\rho.$$

On the complement, use the bound 2 and

$$\sum_{\gamma \notin \Gamma} |\widehat{g}(\gamma)|^4 \leq \lambda^2 M^2 N^2 \sum_{\gamma} |\widehat{g}(\gamma)|^2 \leq \alpha \lambda^2 M^4 N^4.$$

Fourier inversion, with its factor $1/N$, now yields

$$|h(x+t) - h(x)| \leq 2\pi\rho\xi L + 2\alpha\lambda^2 M^4 N^3 = \left(2\pi\rho\xi + \frac{2\lambda^2}{\alpha^2}\right) L.$$

This proves all the claims. □

II.38.1 The precise improvement over the original estimate

The original large-spectrum estimate bounds the sum of fourth powers on Γ by $d(\alpha MN)^4$, and then bounds d by α/λ^2 . The available moment estimate (II.38.2) is stronger by a factor α^2/λ^2 . To achieve the same total translation error $4\alpha\lambda^2 M^4 N^3$, one may therefore replace

$$\rho_{\text{old}} = \frac{\lambda^4}{\pi\alpha^4} \quad \text{by} \quad \rho_{\text{new}} = \frac{\lambda^2}{\pi\alpha^2}.$$

With $\lambda = 2^{-37}\alpha^{11/2}$, this is the change

$$\frac{2^{-148}\alpha^{18}}{\pi} \longrightarrow \frac{2^{-74}\alpha^9}{\pi}. \tag{II.38.4}$$

No character has been added or removed from Γ . The gain comes entirely from a better estimate of the same Fourier expression.

Proposition II.38.2 (Optimal allocation within the moment bounds). *Fix a desired relative translation error $0 < \sigma \leq 1$, with the spectrum and rank defined as in theorem II.38.1, and use only the bounds (II.38.1) and (II.38.3). Among positive choices of λ, ρ satisfying*

$$2\pi\rho + 2\lambda^2/\alpha^2 \leq \sigma,$$

the lower bound for ρ/d obtained by substituting $d \leq \alpha/\lambda^2$ is maximized at

$$\lambda^2 = \frac{\sigma\alpha^2}{4}, \quad \rho = \frac{\sigma}{4\pi}.$$

Its maximum is $\sigma^2\alpha/(16\pi)$. After the 3d radius-selection step, the resulting guaranteed inner radius is $\sigma^2\alpha/(96\pi)$.

Proof. Write $\lambda^2 = q\sigma\alpha^2$, $0 < q < 1/2$. Saturating the error budget gives $\rho = \sigma(1 - 2q)/(2\pi)$, so

$$\frac{\rho}{d} \geq \frac{\rho\lambda^2}{\alpha} = \frac{q(1 - 2q)\sigma^2\alpha}{2\pi}.$$

The quadratic $q(1 - 2q) = 1/8 - 2(q - 1/4)^2$ is maximized at $q = 1/4$. Divide by 6 for the selected inner radius. $\square$

This is an optimization of the guaranteed expression furnished by these estimates. It does not claim optimality of the actual radius for each input function. If the measured fourth moment is smaller, the factor ξ yields the further admissible choice $\rho = \sigma/(4\pi\xi)$, provided it is within the range used for radius selection.

II.39 Synchronizing a local model across fibres

Source 12, delivered Section 7.

The extension theorem concerns a map on a set of directions. Gowers's argument also requires it to agree with differences between elements of another, possibly multivalued, domain. The next lemma makes this second use of overlap explicit.

Let Ω be a finite set with an index map $r : \Omega \rightarrow G$ and a value map $F : \Omega \rightarrow H$. Write $\Omega_a = r^{-1}(a)$ for its fibres. Let $S \subseteq B \subseteq G$, $|S| = a|B|$, and let $\psi : S \rightarrow H$ be a Freiman homomorphism with canonical difference map D .

Lemma II.39.1 (Common-fibre synchronization). *Let $W_1 \subseteq \Omega$. Suppose that for each $w \in W_1$ there is a set $T_w \subseteq S$ with $|T_w| \geq (1 - q)|S|$, and for each $d \in T_w$ the fibre $\Omega_{r(w)+d}$ is nonempty and*

$$|\{z \in \Omega_{r(w)+d} : F(z) - F(w) = \psi(d)\}| \geq (1 - \theta)|\Omega_{r(w)+d}|.$$

For $w_1, w_2 \in W_1$, if $c = r(w_1) - r(w_2)$ satisfies $b_B(c) \leq \kappa$ and

$$2(1 - q)a > 1 + \kappa, \quad 2\theta < 1, \tag{II.39.1}$$

then $c \in S - S$ and

$$F(w_1) - F(w_2) = D(c).$$

Proof. Both $T_{w_1} \subseteq B$ and $T_{w_2} - c \subseteq B - c$. The two-set form of lemma II.34.3, or direct subtraction of their complements, gives

$$|T_{w_1} \cap (T_{w_2} - c)| \geq [2(1 - q)a - 1 - \kappa]|B| > 0.$$

Choose d in this intersection, so $d \in T_{w_1}$ and $d + c \in T_{w_2}$. Their fibres coincide because

$$r(w_1) + d = r(w_2) + (d + c).$$

This common fibre is nonempty. The two subsets on which the required value equations hold have combined missing fraction at most $2\theta < 1$, so they contain a common z . Subtracting the two equations gives

$$F(w_1) - F(w_2) = \psi(d + c) - \psi(d) = D(c).$$

The same choice of d also establishes $c \in S - S$. $\square$

[write] **6 October 2026; sources 28–33.** Source 30's Theorem II.133.1 is this lemma, with the same condition $2(1 - \sigma)a_0 > 1 + \kappa$, proved again without citation.

If $a \geq 1 - \varepsilon$, the convenient sufficient condition is $2(1 - q)(1 - \varepsilon) > 1 + \kappa$. This retains the actual relative density and improves the coarse overlap bookkeeping used in the source's Lemma 10.12.

II.40 A strengthened version of the Section 10 local-model theorem

Source 12, delivered Section 8.

We now specialize to $G = \mathbb{Z}/N\mathbb{Z}$, $N \geq 2$. Primality is unnecessary until the later scalar-linearity corollary. A *multifunction domain* is a finite set X with an index map $r : X \rightarrow G$, or equivalently a partition $X = \bigsqcup_{s \in G} X_s$. A function $\phi : X \rightarrow G$ is a $(1 - \eta)$ -homomorphism of order j if at least a $1 - \eta$ fraction of all ordered tuples

$$(x_1, \dots, x_{2j}) \in X^{2j}, \quad \sum_{i=1}^j r(x_i) = \sum_{i=j+1}^{2j} r(x_i),$$

also satisfy the analogous equality of their ϕ -values. The denominator counts all such tuples, including repeated elements.

II.40.1 The established component-extraction input

Fix, for the remainder of this section,

$$\eta = 2^{-43}, \quad \theta = 10\eta^{1/5}, \quad q = \sqrt{\theta}. \quad (\text{II.40.1})$$

One has

$$\theta < 1/36, \quad q < 1/6. \quad (\text{II.40.2})$$

Indeed, $\theta < 1/36$ is equivalent to $360^5 < 2^{43}$, an integer inequality.

Lemma II.40.1 (Gowers's component input, with explicit hypotheses). *Let $\Omega = \bigsqcup_{s \in G} \Omega_s$, with $|\Omega| = \alpha LN$, $|\Omega_s| \leq L$, where L is a positive integer and $0 < \alpha \leq 1$. Let $F : \Omega \rightarrow G$ be a $(1 - \eta)$ -homomorphism of order two. Suppose $B = -B$ is nonempty and*

$$||\Omega_{s+b}| - |\Omega_s|| \leq \sigma L \quad (s \in G, b \in B).$$

Let $\varrho > 0$ satisfy

$$\varrho \leq \min\{\alpha/192, \alpha^2/32\}, \quad (\text{II.40.3})$$

$$0 < \sigma \leq \min\{\eta\alpha^2, \eta\varrho^4\alpha/384, \eta\varrho\alpha^2/16\}, \quad \sigma < \alpha^2/16. \quad (\text{II.40.4})$$

Then there exist $W_1 \subseteq W \subseteq \Omega$, $S \subseteq B$, and a Freiman homomorphism $\psi : S \rightarrow G$, such that

$$|W| \geq \varrho^2 \alpha^2 LN/16, \quad (\text{II.40.5})$$

$$|W_1| \geq (1 - q)|W|, \quad |S| \geq (1 - \theta)|B|, \quad (\text{II.40.6})$$

and the hypotheses about the sets T_w and the good fibre fractions in lemma II.39.1 hold with this W_1 , S , q , and θ .

Source and deduction. This is the required output of Gowers's Lemmas 10.3–10.9, in the corrected form recorded in [64, 65]. Lemma 10.3 selects an anchor under $\sigma \leq \eta\alpha^2$. Lemma 10.5 selects a fibre-saturated component W under (II.40.3) and $\sigma \leq \eta\varrho^4\alpha/384$, with (II.40.5) and fibre sizes at least $\alpha^2L/16$ on W . Lemma 10.6 supplies S and the nested $1 - \theta$ difference model under the remaining extraction bound. Lemma 10.7 reverses the two outer counting quantifiers, losing $q = \sqrt{\theta}$ and defining W_1 . Lemmas 10.8–10.9 give the exact Freiman property using $6q < 1$ and $W \neq \emptyset$.

The fibre sizes needed in lemma II.39.1 are positive: for $w \in W$ and $d \in S \subseteq B$,

$$|\Omega_{r(w)+d}| \geq |\Omega_{r(w)}| - \sigma L \geq (\alpha^2/16 - \sigma)L > 0.$$

This last point is a necessary hypothesis of the synchronization argument, not a consequence of a proportionate statement on a possibly empty fibre.

The printed proof of Lemma 10.6 has an ambiguity involving a set depending on the direction and a later exchange of quantifiers. The corrected proof used here works on a fixed product of a low-error set with B and chooses each direction's value from an anchor-fibre mode. Appendix II.42 records the count behind this repair. No Bohr-specific property of B is used in this component input. $\square$

II.40.2 The parameterized theorem

Theorem II.40.2 (Local model with a polynomial radius). *Let $X = \bigsqcup_{s \in \mathbb{Z}/N\mathbb{Z}} X_s$ be finite, $M \geq 1$ an integer, $|X_s| \leq M$, and $|X| = \alpha MN$ with $0 < \alpha \leq 1$. Suppose $\phi : X \rightarrow \mathbb{Z}/N\mathbb{Z}$ is a $(1 - 2^{-43})$ -homomorphism of order eight. Choose ϱ, σ satisfying (II.40.3)–(II.40.4), let $\lambda > 0$, and put*

$$\lambda^2 = \frac{\sigma\alpha^2}{4}, \quad \rho = \frac{\sigma}{4\pi}, \quad \Gamma = \{\gamma : |\widehat{g}(\gamma)| \geq \lambda MN\}, \quad g(s) = |X_s|. \quad (\text{II.40.7})$$

Then $1 \leq d = |\Gamma| \leq 4/(\sigma\alpha)$. On

$$C = \mathcal{B}(\Gamma, \zeta), \quad \zeta = \frac{\sigma}{24\pi d} \geq \frac{\sigma^2\alpha}{96\pi}, \quad (\text{II.40.8})$$

there is a Freiman homomorphism $\psi_1 : C \rightarrow \mathbb{Z}/N\mathbb{Z}$ with $\psi_1(0) = 0$, and there is $Y \subseteq X$ with

$$|Y| \geq \frac{(1-q)\varrho^2\alpha}{16}|X|, \quad (\text{II.40.9})$$

such that, whenever $y, z \in Y$ and $r(y) - r(z) \in C$,

$$\phi(y) - \phi(z) = \psi_1(r(y) - r(z)). \quad (\text{II.40.10})$$

Proof. We give the full parameter propagation.

1. The fourfold domain. Let $\Omega = X^4$, with

$$r_\Omega(x_1, x_2, x_3, x_4) = r(x_1) + r(x_2) - r(x_3) - r(x_4)$$

and the identical signed expression defining F from ϕ . Set

$$L = M|X|^3 = \alpha^3 M^4 N^3.$$

Each fibre of Ω has at most $M|X|^3 = L$ elements: after choosing three entries, the index of the fourth is determined, and at most M elements have that index. Moreover

$$|\Omega| = |X|^4 = \alpha LN. \quad (\text{II.40.11})$$

An additive quadruple in Ω expands, after a permutation of its 16 entries, into an equality between two sums of eight indices in X . This is a bijective rearrangement of all tested tuples, and the same rearrangement applies to values. Thus F is a $(1 - \eta)$ -homomorphism of order two.

2. Fourier control and nonempty spectrum. The fibre-count function of Ω is $h = g * \tilde{g} * g * \tilde{g}$. By theorem II.38.1 and (II.40.7), it is $(\mathcal{B}(\Gamma, \rho), \sigma L)$ -invariant. The rank bound is $d \leq 4/(\sigma\alpha)$. The parameter conditions imply $\sigma < 1$, hence $\lambda < \alpha$, so the zero character belongs to Γ and $d \geq 1$. They also imply $0 < \rho < 1/2$.

3. Choose the ambient Bohr set before extraction. Apply theorem II.37.2 with $m = 3d$. It provides $r_0 \in (\rho/2, \rho]$ such that $B = \mathcal{B}(\Gamma, r_0)$ satisfies

$$b_B(c) \leq \kappa := 1 - 4^{-1/3} < 3/8 \quad (c \in C = \mathcal{B}(\Gamma, \rho/(6d))).$$

Since $B \subseteq \mathcal{B}(\Gamma, \rho)$, the already established invariance holds on B . Now, and only now, apply lemma II.40.1 to obtain fresh W, W_1, S, ψ . The character set is unchanged.

4. Extend the directional model. The actual missing density of S is at most $\theta < 1/36$. Therefore

$$3\theta + 2\kappa < \frac{1}{12} + \frac{3}{4} = \frac{5}{6} < 1.$$

By theorem II.35.2, the canonical difference map of ψ induces a Freiman homomorphism ψ_1 on C , with zero value at zero.

5. Synchronize values on the selected component. For the parameters of lemma II.39.1, $a \geq 1 - \theta > 35/36$, $1 - q > 5/6$, and $1 + \kappa < 11/8$. Thus

$$2(1 - q)a > 2 \cdot \frac{5}{6} \cdot \frac{35}{36} = \frac{175}{108} > \frac{11}{8} > 1 + \kappa,$$

and $2\theta < 1$. The synchronization lemma gives

$$F(w_1) - F(w_2) = \psi_1(r_\Omega(w_1) - r_\Omega(w_2))$$

for $w_1, w_2 \in W_1$ with index difference in C .

6. Retain the full slice count. By (II.40.5)–(II.40.11),

$$|W_1| \geq (1 - q) \frac{\varrho^2 \alpha^2 LN}{16} = \frac{(1 - q) \varrho^2 \alpha}{16} |X|^4.$$

Average over the last three coordinates. Some $(x_2, x_3, x_4) \in X^3$ leaves at least the fraction in (II.40.9) of X in the fibre

$$Y = \{y \in X : (y, x_2, x_3, x_4) \in W_1\}.$$

On this slice, differences of indices and of values are exactly $r(y) - r(z)$ and $\phi(y) - \phi(z)$, proving (II.40.10). Finally, the bound on d gives (II.40.8). $\square$

[write] **6 October 2026; sources 28–33.** Source 28's Theorem F.33.1 replaces the component selection by logarithmic windows: retained fraction $\alpha p(\alpha)/20$, $O(\alpha^{-3})$ original-domain frequencies, radius $\Omega(\alpha^5)$, input defect 2^{-21} (Section II.121); it credits this theorem's α^5 retention, extension and smoothing.

II.40.3 Keeping exactly the original large spectrum

Corollary II.40.3 (Explicit sparse-density specialization). *Under the hypotheses on X, ϕ in theorem II.40.2, suppose $0 < \alpha \leq 1/6$. Keep Gowers's threshold*

$$\lambda = 2^{-37}\alpha^{11/2}, \quad \Gamma = \{\gamma : |\widehat{g}(\gamma)| \geq \lambda MN\}.$$

Then $d = |\Gamma| \leq 2^{74}\alpha^{-10}$, and one can take

$$C = \mathcal{B}\left(\Gamma, \frac{2^{-74}\alpha^9}{6\pi d}\right), \quad |Y| \geq \frac{\alpha^5}{20000}|X|. \quad (\text{II.40.12})$$

The local model identity (II.40.10) holds on C , whose radius is at least $\alpha^{19}/(3\pi 2^{149})$.

Proof. Choose $\varrho = \alpha^2/32$ and $\sigma = 2^{-72}\alpha^9$. The bound $\varrho \leq \alpha/192$ is exactly $\alpha \leq 1/6$. The most restrictive smallness comparison is

$$\sigma = 2^{-72}\alpha^9 \leq \frac{2^{-70}}{3}\alpha^9 = \frac{\eta\varrho^4\alpha}{384}.$$

The other conditions in (II.40.4) follow directly from $\alpha \leq 1$; for example $\sigma \leq \eta\varrho\alpha^2/16$ reduces to $\alpha^5 \leq 2^{20}$. Then (II.40.7) gives exactly the displayed λ . Also

$$\frac{(1-q)\varrho^2\alpha}{16} > \frac{5\alpha^5}{98304} \geq \frac{\alpha^5}{20000}.$$

Substitute the rank bound into the radius and simplify. □

[write] **6 October 2026; sources 11–14.** *Against the formal statement* (HEAD e1db00e99). `theorem_10_13` (Section10.lean, line 473; exact companion `theorem_10_13_holds`, Proofs10Main.lean, line 1292) concludes $|Y| \geq \alpha^6|X|/20000$ on the Bohr set of radius `section10Zeta`, under $0 < \alpha \leq 1/6$. The fraction $\alpha^5/20000$ here is below the $5\alpha^5/98304$ that the formal proof already establishes internally (note in Section F.17), so the improvement in the retained fraction is a proof-level fact of the development, not new; the polynomial radius is new. Not formalized.

II.40.4 An all-density schedule and a moment-sensitive version

Corollary II.40.4 (All densities). *For every $0 < \alpha \leq 1$, the choices*

$$\varrho = \min\{\alpha/192, \alpha^2/32\}, \quad \sigma = \frac{2^{-43}\alpha\varrho^4}{384} \quad (\text{II.40.13})$$

satisfy the conditions of theorem II.40.2. Thus its explicit radius and retained-fraction conclusions hold without the restriction $\alpha \leq 1/6$. [6 October 2026; sources 11–14] This schedule is Gowers's own printed choice (transcription `sz-t hm-gower s-pro of .tex`, line 2201: $\rho = \min\{\alpha/192, \alpha^2/32\}$, $\sigma = \eta\rho^4\alpha/384$); it is new relative to the formal statement, which assumes $\alpha \leq 1/6$.

On the branch $\alpha \leq 1/6$, these choices give

$$\sigma = \frac{2^{-70}}{3}\alpha^9, \quad \lambda^2 = \frac{2^{-72}}{3}\alpha^{11}, \quad d \leq 3 \cdot 2^{72}\alpha^{-10}, \quad \zeta \geq \frac{2^{-145}}{27\pi}\alpha^{19}.$$

This last constant is 16/9 times the radius lower bound in corollary II.40.3, but uses a slightly different spectral threshold.

Proof. The only remaining comparisons are $\varrho^4 \leq 384\alpha$, $\varrho^3 \leq 24\alpha$, and $\eta\varrho^4 < 24\alpha$. They follow from $\varrho \leq \alpha/192$, $0 < \alpha \leq 1$, and $\eta < 1$. These inequalities give the three conditions on σ and fibre positivity. The sparse-branch formulas follow by substituting $\varrho = \alpha^2/32$ and simplifying powers of 2 and 3. $\square$

Corollary II.40.5 (Using the measured fourth moment). *In theorem II.40.2, let $\xi = (LN)^{-1} \sum |\hat{g}|^4$. If $\rho_* := \sigma/(4\pi\xi) \leq 1/2$, the same spectrum and retained fraction work on the larger Bohr set of radius*

$$\zeta_* = \frac{\sigma}{24\pi\xi d} \geq \frac{\sigma^2\alpha}{96\pi\xi}.$$

For the schedule (II.40.13), the required bound on ρ_ is automatic.*

Proof. Use the precise version of (II.38.3). Nothing after the smoothing step changes. For the final assertion, $\xi \geq \alpha$ gives $\rho_* \leq \eta\varrho^4/(1536\pi) < 1/2$. $\square$

II.41 Consequences and the size of the improvement

Source 12, delivered Section 9.

II.41.1 The volume cost changes from quadratic to linear in the rank

Let $k = \lceil 2^{74}\alpha^{-10} \rceil$ and compare the explicit rank-based radius

$$\zeta_{\text{old}} = 2^{-155k} \alpha^{18k}/k \quad \text{with} \quad \zeta_{\text{new}} = \frac{\alpha^{19}}{3\pi 2^{149}}.$$

The latter is a guaranteed smaller radius inside the actual set in corollary II.40.3; hence the local model also holds on $\mathcal{B}(\Gamma, \zeta_{\text{new}})$. The size bound (II.37.1) gives

$$\begin{aligned} |\mathcal{B}(\Gamma, \zeta_{\text{old}})| &\geq 2^{-E_{\text{old}}} N, \\ E_{\text{old}} &= k[1 + 155k + 18k \log_2(1/\alpha) + \log_2 k], \\ |\mathcal{B}(\Gamma, \zeta_{\text{new}})| &\geq 2^{-E_{\text{new}}} N, \\ E_{\text{new}} &= k[150 + \log_2(3\pi) + 19 \log_2(1/\alpha)]. \end{aligned}$$

Here using k rather than the actual d only weakens a lower bound because the base is below one. Thus the certified logarithmic volume cost changes from

$$O(k^2[1 + \log(1/\alpha)]) \quad \text{to} \quad O(k[1 + \log(1/\alpha)]).$$

The large rank itself remains. This is a substantial improvement at this stage of the proof, but the resulting numerical constants are still enormous.

Example II.41.1 (Density $\alpha = 1/8$). Then $k = 2^{104}$. The old negative base-two logarithm of the radius is

$$\log_2(1/\zeta_{\text{old}}) = 209 \cdot 2^{104} + 104.$$

The new guaranteed radius has

$$\log_2(1/\zeta_{\text{new}}) = 206 + \log_2(3\pi) \approx 209.2365.$$

The retained fraction improves by a factor of 8 relative to the corrected $\alpha^6/20000$ statement. [6 October 2026; sources 11–14] Relative to what the formal proof already establishes ($5\alpha^5/98304$, note in Section F.17) there is no gain in the fraction; the gain is in the radius. These are parameter comparisons; they do not assert that a set of feasible computational size meets the full theorem’s worst-case volume bound.

II.41.2 Long progressions of locally linear directions

Corollary II.41.2 (Scalar linearity along short Bohr progressions). *Let $N = p$ be prime, and suppose $Y, \psi_1, C = \mathcal{B}(\Gamma, \zeta)$ satisfy theorem II.40.2. If $m \geq 1$, $0 \neq h \in \mathcal{B}(\Gamma, \zeta/m)$, and $m < p$, then there is $a \in \mathbb{Z}/p\mathbb{Z}$ such that*

$$\phi(y) - \phi(z) = a(r(y) - r(z))$$

whenever $y, z \in Y$ and $r(y) - r(z) \in \{jh : -m \leq j \leq m\}$.

Proof. All jh , $|j| \leq m$, belong to C by the triangle inequality. Since ψ_1 is a Freiman homomorphism and $\psi_1(0) = 0$, the relations

$$jh + h = (j + 1)h + 0$$

give $\psi_1(jh) = j\psi_1(h)$ by induction for nonnegative j , and oddness gives the negative values. Take $a = \psi_1(h)/h$, which is legitimate since $h \neq 0$ in the prime field. Then use (II.40.10). $\square$

Corollary II.41.3 (An explicit supply of nonzero directions). *If $0 < \zeta \leq 1/2$ and $\zeta p^{1/d} \geq 4$, put*

$$m = \left\lfloor \zeta p^{1/d} / 4 \right\rfloor.$$

Then $\mathcal{B}(\Gamma, \zeta/m)$ contains a nonzero h , and the progression $P = \{jh : -m \leq j \leq m\} \subseteq C$ is proper. A single scalar governs all index differences of pairs from Y which belong to P .

Proof. We have $m \geq 1$ and

$$|\mathcal{B}(\Gamma, \zeta/m)| \geq \left(\frac{\zeta}{2m} \right)^d p \geq 2^d > 1.$$

The elementary size bound remains valid if $\zeta/m = 1/2$. Also $m \leq p/8$, since $d \geq 1$ and $\zeta \leq 1/2$. Thus $2m < p$, so the distinct integers $-m, \dots, m$ give distinct multiples of a nonzero element in $\mathbb{Z}/p\mathbb{Z}$. Apply corollary II.41.2. $\square$

This supplies the direct analogue of Gowers’s Corollary 10.14 with the larger radius. It does not assert that Y itself contains a progression of this length; it asserts that differences in the displayed progression are governed by one scalar.

II.41.3 What does not yet follow

The later proof still requires compatible multilinear structure, partitions with controlled widths, recurrence thresholds, and an iteration which pays for the density lost when changing cyclic models. The current repository records unresolved obligations in Sections 13, 16, and 18 [66]. [6 October 2026; sources 11–14] *Incomplete list:* the ledger at the pin also had open statements in Sections 1, 5, 7 and 8, and at HEAD e1db00e99 the 20 open statements are `theorem_1_2`, `theorem_1_3`, `corollary_5_8`, `theorem_7_1`, `theorem_7_2`, `theorem_8_2`, `corollary_8_3`, `lemma_13_5`–`lemma_13_7`, `lemma_13_9`, `corollary_13_10`, `theorem_13_12`, `theorem_16_2`, `lemma_16_4`,

lemma_16_10, corollary_16_11, theorem_18_1, theorem_18_2 and corollary_18_7. A local radius improvement cannot be substituted into every later theorem without checking these dependencies.

In particular, the theorem established here is a complete replacement of the Section 10 local-model output under its stated assumptions. It does not prove a better final $r_\ell(N)$ bound. Determining how much of the local gain survives the entire argument is one of the proposed research projects below.

II.42 The fixed-domain repair in the Section 10 extraction input

Source 12, delivered Appendix A.

This appendix explains the specific quantifier issue checked when using lemma II.40.1. It is not a reproof of every component-selection lemma. The construction below is the one present in the exact companion proof of Lemma 10.6 in `Proofs10Extraction.lean` [65].

Let $t = \eta^{1/5}$, so $0 < t < 1/10$ for the value of η used here. Let x be the anchor selected by Lemma 10.3, and let W be the component selected by Lemma 10.5. For two domain points x, w , let $b(x, w)$ count pairs (y, z) with

$$r(y) - r(x) = r(z) - r(w) \in B,$$

and let $e(x, w)$ count those pairs for which

$$F(y) - F(x) \neq F(z) - F(w).$$

Write $\epsilon(x, w) = e(x, w)/b(x, w)$ when the denominator is positive. The hypotheses below ensure that it is positive for $w \in W$.

Define the *fixed* low-error set

$$\mathcal{L} = \{w \in W : \epsilon(x, w) \leq 300t^4\}.$$

The regular-component estimate gives $|\mathcal{L}| \geq (1 - 5t)|W|$. For $w \in \mathcal{L}$ and $d \in B$, set

$$A_d = \Omega_{r(x)+d}, \quad C_{w,d} = \Omega_{r(w)+d}, \quad P_w(d) = |A_d| |C_{w,d}|.$$

The anchor and component fibre bounds, together with pointwise invariance and the inherited bound on ϱ , imply that all these fibres are nonempty and that

$$P_w(d) \leq 4P_w(d') \quad (d, d' \in B).$$

For example, it suffices that each translated fibre size lies between $3/4$ and $5/4$ of its unshifted size; the product ratio is then at most $(5/3)^2 < 4$. The parameter schedule of theorem II.40.2 is stronger than this requirement.

Disintegrating over the *fixed set of directions* B gives exact identities

$$b(x, w) = \sum_{d \in B} P_w(d), \quad e(x, w) = \sum_{d \in B} \#\{\text{disagreeing pairs in } A_d \times C_{w,d}\}.$$

Let $p = \min_{d \in B} P_w(d) > 0$. Then

$$e(x, w) \leq 300t^4 b(x, w) \leq 1200t^4 |B| p.$$

If more than $35t^2|B|$ directions each had more than a $35t^2$ fraction of disagreeing pairs, the last count would exceed

$$35^2t^4|B|p = 1225t^4|B|p,$$

a contradiction. Thus, for each $w \in \mathcal{L}$, at most a $35t^2$ fraction of directions are bad.

Now double-count the bad pairs in $\mathcal{L} \times B$. At most a $7t$ fraction of directions can have more than a $5t$ fraction of $\mathcal{L}$ bad, since $(7t)(5t) = 35t^2$. Retain the remaining set S of directions. It has size at least $(1 - 7t)|B|$, hence at least $(1 - 10t)|B|$. For each retained direction, the good points of $\mathcal{L}$ comprise at least

$$(1 - 5t)|\mathcal{L}| \geq (1 - 5t)^2|W| \geq (1 - 10t)|W|.$$

For each $d \in S$, choose $\psi(d)$ to be the mode of $F(y) - F(x)$ on the anchor fibre A_d . This choice depends on d but is independent of w . For any good pair (w, d) , pair agreement at least $1 - 35t^2$ forces that mode to occupy at least a $1 - 35t^2$ fraction of A_d . Every product of a point in the mode class with an incorrectly labelled point of $C_{w,d}$ is a disagreeing pair. Hence the incorrectly labelled fraction of $C_{w,d}$ is at most

$$\frac{35t^2}{1 - 35t^2} \leq 10t.$$

This proves the nested $1 - 10t$ conclusion on the fixed W .

The proof never exchanges quantifiers over a set of the form $W \cap (W - d)$ treated as if it were independent of d . Its invariance input is pointwise for each $d \in B$, which is why choosing the smaller ambient Bohr radius before component extraction is legitimate. The later exact Freiman step, Lemma 10.9, likewise uses a fixed W and the strict bound $6\sqrt{\theta} < 1$.

II.43 Sharp stability of the cross-section collision bound

Source 14, delivered Section 4.

The arrangement argument in Sections 12 and 14 of Gowers's proof passes through a second moment that counts pairs of congruent cubes in the same cross-section. A useful profile bound is

$$R_1 \leq \sum_r |B_r|^3.$$

The numerical inequality conceals a strong geometric restriction: a cross-section can nearly attain its cubic upper bound only when it is nearly a product of subgroup cosets. We prove a dimension-independent stability statement, classify exact equality including the labeling, and explain the resulting extra information in the arrangement moment inequalities. The group-theoretic ingredient is the classical three-halves difference-set lemma, for which we include the short proof. Our sharpness assertion concerns an infinite family of exact equality examples and the leading asymptotic constant; we do not assert that every real value of the error parameter is attained by a finite set.

II.43.1 Cube collisions and two elementary upper bounds

Let $G_1, \dots, G_k$ be finite abelian groups, where $k \geq 1$, and put $G = \prod_{i=1}^k G_i$. Write $e_i : G \rightarrow G$ for the coordinate projection followed by the coordinate inclusion: $e_i(x)$ has coordinate x_i in position i and zero elsewhere. If $x, h \in G$, the axis-parallel cube with base x and sidelengths h has vertices

$$x + \omega \cdot h = (x_i + \omega_i h_i)_{i=1}^k, \quad \omega \in \{0, 1\}^k.$$

Zero sidelengths and repeated vertices are included throughout.

Let $A \subset G$, and let $\phi : A \rightarrow K$, where K is an abelian group. For a cube contained in A define

$$D_h\phi(x) = \sum_{\omega \in \{0,1\}^k} (-1)^{k-|\omega|} \phi(x + \omega \cdot h).$$

Let $Q_\phi(A)$ count triples $(x, z, h) \in G^3$ for which both cubes are contained in A and $D_h\phi(x) = D_h\phi(z)$. Let $Q(A)$ denote the same count without the value condition. With $m = |A|$, put

$$E(A) = \#\{(x, z, w) \in A^3 : z + w - x \in A\}.$$

The injective change of variables $(x, z, h) \mapsto (x, z, x + h)$ gives

$$Q_\phi(A) \leq Q(A) \leq E(A) \leq m^3. \quad (\text{II.43.1})$$

For every coordinate i , also define

$$C_i(A) = \#\{(x, w) \in A^2 : x + e_i(w - x) \in A\}.$$

The same change of variables shows that

$$Q_\phi(A) \leq mC_i(A). \quad (\text{II.43.2})$$

Indeed, a failed corner condition for (x, w) excludes every possible choice of the second base z .

II.43.2 A sharp near-maximal energy lemma

We first prove a statement that applies to an arbitrary finite abelian group, without a specified product structure. Define

$$\tau(\varepsilon) = \frac{1 - \sqrt{1 - 4\varepsilon}}{2} = \varepsilon + \varepsilon^2 + 2\varepsilon^3 + O(\varepsilon^4).$$

Thus $\tau(\varepsilon)(1 - \tau(\varepsilon)) = \varepsilon$.

Lemma II.43.1 (The three-halves difference-set lemma). *If B is a nonempty finite subset of an abelian group and $|B - B| < 3|B|/2$, then $B - B$ is a subgroup.*

Proof. Put $b = |B|$ and $D = B - B$. If $d = b_1 - b_2 \in D$, the two subsets $B - b_1$ and $B - b_2$ of D have intersection of size at least $2b - |D| > b/2$. Consequently

$$|B \cap (B + d)| > b/2 \quad (d \in D).$$

For $d, e \in D$, the two subsets $B \cap (B + d)$ and $B \cap (B + e)$ of B intersect. Hence $d - e \in B - B = D$. The nonempty finite set D is closed under differences and is a subgroup. $\square$

[write] **6 October 2026; sources 11–14.** This is a classical fact of additive combinatorics (usually credited to Freiman); the source calls it classical and proves it, without a citation.

[write] **6 October 2026; sources 15–20.** Source 15 proves the same lemma (its Lemma II.52.1).

Theorem II.43.2 (Near-maximal energy, with the quadratic error envelope). *Let A be a nonempty subset of a finite abelian group, let $m = |A|$, and let $0 \leq \varepsilon \leq 1/100$. If*

$$E(A) \geq (1 - \varepsilon)m^3,$$

then there is a coset $C = c + H$ of a subgroup such that

$$|A \triangle C| \leq \tau(\varepsilon)m. \quad (\text{II.43.3})$$

The coset satisfying this bound is unique, and is the unique coset minimizing symmetric difference with A .

Proof. Write $r(d) = |A \cap (A + d)|$ and $\mathcal{D} = m^3 - E(A)$. The identities

$$\sum_d r(d) = m^2, \quad E(A) = \sum_d r(d)^2$$

give

$$\sum_{a_0, b \in A} (m - r(b - a_0)) = \mathcal{D}.$$

Choose $a_0 \in A$ such that

$$\sum_{b \in A} (m - r(b - a_0)) \leq \varepsilon m^2$$

and define

$$B = \{b \in A : r(b - a_0) \geq 9m/10\}.$$

Then $|B| \geq (1 - 10\varepsilon)m \geq 9m/10$. The elementary overlap inequality

$$r(d - e) \geq r(d) + r(e) - m$$

shows that $r(d) \geq 4m/5$ for every $d \in B - B$. Therefore

$$|B - B| \leq \frac{5m}{4} < \frac{3}{2}|B|.$$

By Lemma II.43.1, $H = B - B$ is a subgroup. Choose $b_0 \in B$ and let $C = b_0 + H$. If

$$A_0 = A \cap C, \quad a = |A_0|, \quad t = |A \setminus C|, \quad u = |C \setminus A|, \quad h = |H|,$$

then

$$a \geq 9m/10, \quad t \leq m/10, \quad h \leq 5m/4. \quad (\text{II.43.4})$$

We now sharpen this rough approximation by counting disjoint families of triples in A^3 whose alternating sum is outside A . For every $v \in C \setminus A$, and every middle entry $y \in A_0$, the number of pairs $x, z \in A_0$ with $x - y + z = v$ is at least $2a - h$. This follows by intersecting A_0 and $(v + y) - A_0$, two subsets of the coset C , both of size a . These triples contribute at least $ua(2a - h)$ to $\mathcal{D}$.

Next consider triples with exactly one entry in $A \setminus C$. There are three positions for that entry. Fix one of its t values and one of the three positions. The other two entries range over A_0^2 . Their alternating sum is outside C . For it to belong to A , it must therefore belong to the t -element set $A \setminus C$. Fixing one of the two entries in A_0 and this output determines the remaining entry, so at

most *at* of the a^2 choices give an output in A . At least $a(a-t)$ choices fail. These three families are disjoint from one another and from the all-inside family. We obtain

$$\mathcal{D} \geq ua(2a-h) + 3ta(a-t). \quad (\text{II.43.5})$$

Using (II.43.4), the coefficients on the right are at least

$$a(2a-h) \geq \frac{99}{200}m^2, \quad 3a(a-t) \geq \frac{54}{25}m^2.$$

In particular, if $\delta = (t+u)/m$ then

$$\delta \leq \frac{200}{99}\varepsilon < \frac{1}{3}.$$

Put $x = t/m$; thus $0 \leq x \leq \delta$, $a/m = 1-x$, and $h/m = 1+\delta-2x$. Dividing (II.43.5) by m^3 and simplifying gives

$$\begin{aligned} \frac{\mathcal{D}}{m^3} &\geq (1-x)[(\delta-x)(1-\delta) + 3x(1-2x)] \\ &= \delta(1-\delta) + x[2+\delta^2 - (8+\delta)x + 6x^2]. \end{aligned}$$

For $0 \leq x \leq \delta \leq 1/3$, the bracket is nonnegative. Indeed, its derivative in x is $-(8+\delta) + 12x < 0$ on this interval, and its value at $x = \delta$ is $2(1-\delta)(1-3\delta) \geq 0$. Thus

$$\delta(1-\delta) \leq \frac{\mathcal{D}}{m^3} \leq \varepsilon.$$

Since $\delta < 1/3$, inversion of this increasing quadratic proves (II.43.3).

For uniqueness, suppose two cosets C, C' satisfy $|A\Delta C|, |A\Delta C'| \leq \tau m$, with $\tau = \tau(\varepsilon) < 1/5$. Then

$$|C \cap C'| \geq (1-2\tau)m, \quad |C|, |C'| \leq (1+\tau)m.$$

Their intersection is consequently larger than half of each coset. A nonempty intersection of cosets is a coset of the intersection of their subgroups. An index other than one is at least two, so the two subgroups and then the two cosets must coincide. Any closer coset would also satisfy the displayed distance bound, proving the last assertion. $\square$

[write] 7 October 2026; sources 40–46. *Extended by source 45 to $0 \leq \varepsilon < 1/9$ in every abelian group, with the canonical subgroup $\{d : |A \cap (A+d)| > m/2\}$ and all positive equality cases, the punctured cosets (Theorems II.185.4, II.185.8; item II.31).*

[write] 7 October 2026; sources 60–69. The range $\varepsilon \leq 1/100$ is extended to $1/4$, optimal, by source 68 (Section II.296; item II.31), after source 45's $1/9$.

Remark W3 (Source 11's Question 3 in substance; added at the third write, 6 October 2026). Source 11's set energy $E(S)$ (Section II.23) is this $E(A)$: both count ordered $(x_1, x_2, x_3, x_4) \in A^4$ with $x_1 + x_2 = x_3 + x_4$. Let $f : G \rightarrow W$, an affine ℓ , $S = \{f \neq \ell\}$, $s = |S| = tN$ with $0 < t < 1/7$, W without involutions, and suppose $\varepsilon(f) \leq \Psi_0(t) + \eta$ with $\theta := \eta/(6t^3) \leq 1/100$. Then there are a unique subgroup coset $C = c + H$ of G and a nonzero $b \in W$ such that

$$|S\Delta C| \leq \tau(\theta)s, \quad |\{x : f(x) - \ell(x) \neq b\mathbf{1}_C(x)\}| \leq \left(\tau(\theta) + \frac{\eta}{2t^2(1-7t)}\right)s.$$

Proof. Source 11's Corollary II.27.1 gives $t^3 - E(S)/N^3 \leq \eta/6$, that is $E(S) \geq (1 - \theta)s^3$; Theorem II.43.2 (applied to the nonempty set S) gives the unique coset C with $|S\Delta C| \leq \tau(\theta)s$. Let b be a most frequent nonzero value of $g = f - \ell$ on S and u the number of points of S with another value; the proof of Corollary II.27.1 gives $su \leq A \leq \eta N^2/(2(1 - 7t))$, so $u/s \leq \eta/(2t^2(1 - 7t))$. The map $g - b\mathbf{1}_C$ is nonzero only on $S\Delta C$ and on those u points of $S \cap C$. *With an involution in W* the same argument, using source 11's bounds $t^3 - e_S \leq \eta/8$ and $A/N^2 \leq \eta/(6(1 - 3t))$ for $0 < t < 1/3$, gives $\theta = \eta/(8t^3)$ and $u/s \leq \eta/(6t^2(1 - 3t))$. Both bounds are relative to s and depend on η/t^3 , the scale that source 11's Question 3 asks for; the coset part is sharp to first order by Proposition II.43.4 (with $\phi = 0$, $Q_\phi = E$). Neither source writes this composition (the intake found it); not formalized. What remains of the question: the range $\theta > 1/100$, and whether the value term can be improved.

II.43.3 The product structure is forced, independently of dimension

Theorem II.43.3 (Sharp stability for cube collisions). *Let $A \subset \prod_{i=1}^k G_i$ be nonempty, let $m = |A|$, and let $\phi : A \rightarrow K$ be any function to an abelian group. Suppose*

$$Q_\phi(A) \geq (1 - \varepsilon)m^3, \quad 0 \leq \varepsilon \leq 1/100.$$

Then there are subgroups $H_i \leq G_i$ and points $c_i \in G_i$ such that the unique closest coset to A is the coordinate product

$$C = \prod_{i=1}^k (c_i + H_i), \quad |A\Delta C| \leq \tau(\varepsilon)m. \quad (\text{II.43.6})$$

The statement holds also with $Q(A)$ in place of $Q_\phi(A)$.

Proof. By (II.43.1), Theorem II.43.2 applies. Use the subgroup H , coset $C = c + H$, and notation from its proof. We show that $e_i H \subset H$ for every i .

Suppose instead that $e_i H \not\subset H$. The homomorphism

$$F_i : H \times H \longrightarrow G, \quad F_i(y, w) = y + e_i(w - y),$$

has image $H + e_i H$. To verify this, its image contains H by taking $y = w$, and contains $e_i H$ by taking $y = 0$; the reverse inclusion is immediate. The image has at least $2h$ elements, so every nonempty fiber of F_i has size at most $h/2$. After translation by c , at most $mh/2$ pairs in C^2 have their mixed corner in A .

At most $m^2 - a^2$ pairs in A^2 fail to belong to C^2 . Consequently the rough bounds (II.43.4) yield

$$C_i(A) \leq m^2 - a^2 + \frac{mh}{2} \leq \left(1 - \frac{81}{100} + \frac{5}{8}\right)m^2 = \frac{163}{200}m^2.$$

But (II.43.2) gives $C_i(A) \geq (1 - \varepsilon)m^2 \geq 99m^2/100$, a contradiction. Thus each coordinate component of an element of H is again in H . If H_i is its projection into G_i , it follows that $H = \prod_i H_i$. The distance estimate and uniqueness were already proved in Theorem II.43.2. $\square$

[write] **6 October 2026; sources 15–20.** Source 15's Theorem II.52.2 proves this conclusion, independently and by the same method, with the weaker constants 8ε , 2ε and 18ε ; its Example II.52.3 is Proposition II.43.4. With them the leading constant asked for by source 15's first question is 1 (item II.31).

[write] **7 October 2026; sources 40–46.** Source 45's Theorems II.185.14 and II.185.15 extend the range to defect below $1/9$ and show that positive-defect equality forces one punctured coordinate.

For $G_i = \mathbb{F}_p$, every H_i is either $\{0\}$ or $\mathbb{F}_p$. The approximating set is therefore an affine coordinate subspace: each coordinate is either fixed or completely free. No primality or torsion assumption was needed for the general result.

Proposition II.43.4 (Exact sharpness at infinitely many errors). *For every prime $p \geq 101$ and every $k \geq 1$, let*

$$A = (\mathbb{F}_p \setminus \{0\}) \times \mathbb{F}_p^{k-1}, \quad \phi = 0, \quad \delta_p = \frac{1}{p-1}, \quad \varepsilon_p = \delta_p(1 - \delta_p).$$

Then $\varepsilon_p \leq 1/100$,

$$Q_\phi(A) = (1 - \varepsilon_p)|A|^3, \quad \min_{\text{coordinate cosets } C} \frac{|A \triangle C|}{|A|} = \delta_p = \tau(\varepsilon_p).$$

Thus the envelope is attained at an infinite sequence of errors tending to zero, and its leading coefficient one cannot be decreased.

Proof. Put $n = p - 1$. For $A_1 = \mathbb{F}_p \setminus \{0\}$, the difference multiplicity is n at zero and $n - 1$ at every nonzero element. Therefore

$$E(A_1) = n^2 + n(n - 1)^2 = n^3 - n^2 + n.$$

In the remaining full coordinates, every choice of the two bases and sidelength is valid. Hence

$$Q_\phi(A) = E(A_1)p^{3(k-1)}, \quad |A| = np^{k-1},$$

which proves the count. The full space differs from A in $p^{k-1} = \delta_p|A|$ points. Every proper coordinate coset has at most p^{k-1} points and hence differs from A in at least $(p - 2)p^{k-1} > p^{k-1}$ points. This proves the distance assertion. Theorem II.43.2 in fact shows that the full space is the unique closest coset among all subgroup cosets. $\square$

II.43.4 Exact equality, including the labeling

The value condition in Q_ϕ is sensitive only to the highest interaction among the k coordinates. Lower interactions must be retained when describing equality; declaring the entire labeling multilinear would be false.

Theorem II.43.5 (Complete equality classification for one cross-section). *For a nonempty $A \subset \prod_{i=1}^k G_i$, with $m = |A|$, one has $Q_\phi(A) = m^3$ if and only if both of the following conditions hold.*

1. $A = \prod_i (c_i + H_i)$ for subgroups $H_i \leq G_i$.
2. After translation, $f(x) = \phi(c + x)$ on $\prod_i H_i$ has the form

$$f(x) = M(x_1, \dots, x_k) + \sum_{S \subseteq \{1, \dots, k\}} f_S(x_S), \quad (\text{II.43.7})$$

where M is additive in each variable separately, and the functions $f_S : \prod_{i \in S} H_i \rightarrow K$ are arbitrary. The empty-set term is a constant.

The full interaction M is uniquely determined by f and the chosen origin c , although the lower-term representation need not be unique.

Proof. Theorem II.43.3 with $\varepsilon = 0$ proves the first condition. On this product coset, every triple of the two bases and sidelengths is a valid cube pair. Equality says precisely that $D_h f(x)$ is independent of x for each h . Let

$$M(h) = D_h f(0) = \sum_{S \subseteq \{1, \dots, k\}} (-1)^{k-|S|} f(h_S, 0_{S^c}).$$

The finite-difference identity

$$\Delta_{u+v}^i g(x) = \Delta_u^i g(x) + \Delta_v^i g(x + e_i u)$$

and the base-independence of the full mixed difference imply that M is additive in its i th variable, for every i . Rearranging the defining inclusion-exclusion formula for M gives (II.43.7) with functions of proper subsets of the coordinates. Conversely, the full mixed difference of every proper-subset term is zero, while that of a separately additive M equals $M(h)$ and is independent of the base. This proves sufficiency and also identifies the unique full interaction. $\square$

[write] **6 October 2026; sources 15–20.** Source 15's Theorems II.51.1 and II.51.2 prove the same classification independently, without the uniqueness of the top interaction.

When $k = 1$, the conclusion is an affine homomorphism on a coset. When some H_i is the trivial subgroup, $D_h f$ vanishes identically, so every labeling is allowed; this agrees with (II.43.7), because that coordinate is redundant. For a full cross-section $\mathbb{F}_p^k$ with values in $\mathbb{F}_p$, separate additivity gives

$$M(x) = ax_1 \cdots x_k,$$

but the lower-coordinate functions remain arbitrary. For example, $axy + u(x) + v(y)$ is the complete form in dimension two. Empty cross-sections contribute zero and impose no condition.

II.43.5 Consequences for a family of cross-sections

Let $B \subset G \times R$, where R is a finite abelian group, and let $\phi : B \rightarrow K$. Put $A_r = B_r$, $m_r = |A_r|$, and

$$S = \sum_r m_r^3, \quad R_1 = \sum_r Q_{\phi|_{A_r}}(A_r).$$

Then $R_1 = S$ holds exactly when every nonempty cross-section has the product-coset geometry and labeling described in Theorem II.43.5. The different cross-sections may have different coordinate subgroups and different top interactions.

Suppose $S > 0$ and $R_1 \geq (1 - \varepsilon)S$. For $m_r > 0$ define

$$\varepsilon_r = 1 - \frac{Q_{\phi|_{A_r}}(A_r)}{m_r^3}.$$

Thus $\sum_r \varepsilon_r m_r^3 \leq \varepsilon S$. For any $0 < \theta \leq 1/100$, the cross-sections with $\varepsilon_r \leq \theta$ have total cubic weight at least $(1 - \varepsilon/\theta)S$. Each of them has a unique product coset C_r satisfying

$$|A_r \triangle C_r| \leq \tau(\varepsilon_r) m_r \leq \tau(\theta) m_r.$$

Moreover the edits on these cross-sections satisfy the weighted bound

$$\sum_{\varepsilon_r \leq \theta} m_r^2 |A_r \triangle C_r| \leq \frac{\varepsilon S}{1 - \tau(\theta)}. \quad (\text{II.43.8})$$

Indeed $\tau(u) = u/(1 - \tau(u)) \leq u/(1 - \tau(\theta))$ for $0 \leq u \leq \theta$.

There is also a useful contrapositive. Let

$$d_r = \min_{\text{coordinate cosets } C} \frac{|A_r \Delta C|}{m_r}, \quad \rho_r = \min\{d_r, \tau(1/100)\}$$

for nonempty cross-sections. The stability theorem implies

$$Q_{\phi|A_r}(A_r) \leq m_r^3(1 - \rho_r(1 - \rho_r)). \quad (\text{II.43.9})$$

If the collision deficit is at most $1/100$, invert the stability bound; otherwise the capped right-hand defect is at most $1/100$. Consequently, with

$$\Xi = \sum_{r:m_r>0} m_r^3 \rho_r(1 - \rho_r),$$

one has $R_1 \leq S - \Xi$. Retaining a lower bound for the distance from product cosets therefore strengthens the denominator used by the arrangement interpolation argument.

II.44 Equality and stability in arrangement moment interpolation

Source 14, delivered Section 5.

We recall the exact representation in a form allowing different groups for the cube coordinates, the cross-section index, and the label. Assume now that K is finite, and put $T = R \times K$. For $h \in G$ define a nonnegative integer weight

$$a_h(r, v) = \#\{x \in G : x + \{0, 1\}^k \cdot h \subset B_r, D_h \phi(x, r) = v\}.$$

Let R_j count ordered lists of $2j$ congruent cubes for which the sum of the first j cross-section indices equals the sum of the last j , and likewise for their mixed-difference labels. Character orthogonality gives

$$R_j = \frac{1}{|T|} \sum_{h \in G} \sum_{\chi \in \hat{T}} |\hat{a}_h(\chi)|^{2j}, \quad R_1 = \sum_{h,r,v} a_h(r, v)^2. \quad (\text{II.44.1})$$

The Fourier transform here is unnormalized. Replacing D_h by $(-1)^k D_h$, as in an alternative sign convention in the source, does not change these counts.

Theorem II.44.1 (Complete equality criterion for the moment inequality). *Suppose $R_2 > 0$ and $j \geq 3$. Then*

$$R_j \geq \frac{R_2^{j-1}}{R_1^{j-2}}. \quad (\text{II.44.2})$$

Equality holds if and only if there is a common number $M > 0$ with the following property: for every h , either $a_h = 0$, or there are a subgroup $L_h \leq T$ and a point $z_h \in T$ such that

$$a_h = \frac{M}{|L_h|} 1_{z_h + L_h}. \quad (\text{II.44.3})$$

In particular, all active sidelengths have exactly the same total number M of cubes, whereas their support subgroups may differ. The integer-valued nature of a_h imposes the evident divisibility condition in (II.44.3).

Proof. Set $X_{h,\chi} = |\hat{a}_h(\chi)|^2$ and give this finite index set the probability weights

$$\nu(h, \chi) = \frac{X_{h,\chi}}{|T|R_1}.$$

The mean of X under ν is $\mu = R_2/R_1 > 0$. Jensen's inequality for the strictly convex function $x \mapsto x^{j-1}$ gives (II.44.2); equality is equivalent to every nonzero $X_{h,\chi}$ being the same positive number μ .

For active h , the trivial character gives $\sum_z a_h(z) = \sqrt{\mu} = M$, independently of h . We claim that a nonnegative weight a of mass M whose Fourier coefficients all have modulus either zero or M is a constant multiple of a coset indicator. Choose z_0 in its support and let L be the subgroup generated by all support differences. Equality in the triangle inequality says

$$|\hat{a}(\chi)| = M \iff \chi \text{ is constant on the support of } a \iff \chi \in L^\perp.$$

The assumed zero-or- M alternative thus identifies the entire Fourier transform with that of $(M/|L|)1_{z_0+L}$. Fourier inversion proves the claim. This establishes necessity of (II.44.3). Conversely, each such weight has Fourier modulus M on $L_h^\perp$ and zero elsewhere, proving equality in (II.44.2). $\square$

[write] 6 October 2026; sources 15–20. Source 15's Theorem II.47.3 is the same inequality and equality criterion in a division-free form.

For $j = 2$ the moment inequality is an identity and has no rigidity content. This is why the strict assumption $j \geq 3$ is necessary.

[write] 6 October 2026; sources 11–14. *Credit.* Inequality (II.44.2) with the profile S is source 04's arrangement interpolation (Theorem II.10.1, its (A5) and (A7)), and source 04's equality models (A13), (A14) are special cases of the classification here. This advances source 04's Question 7 (carrying the profile) without answering it (item II.3).

Corollary II.44.2 (Equality in the profile bound for higher arrangements). *If $R_2 > 0$ and $j \geq 3$, then*

$$R_j \geq \frac{R_2^{j-1}}{S^{j-2}}.$$

Equality holds exactly when both the cross-section classification of Theorem II.43.5 and the common-mass coset condition of Theorem II.44.1 hold. More generally,

$$R_j \geq \frac{R_2^{j-1}}{(S - \Xi)^{j-2}}. \quad (\text{II.44.4})$$

The denominator is positive whenever $R_2 > 0$, because $R_1 \leq S - \Xi$ and $R_1 > 0$.

Corollary II.44.3 (A defect computable from cross-section cardinalities). *Suppose the cube coordinates form $\mathbb{F}_p^k$. For $1 \leq m \leq p^k$ define*

$$b_p(m) = \min \left\{ \tau(1/100), \min_{0 \leq a \leq k} \frac{|m - p^a|}{m} \right\}, \quad S_{\text{card}} = \sum_{r: m_r > 0} m_r^3 (1 - b_p(m_r))(1 - b_p(m_r)).$$

Then

$$R_1 \leq S_{\text{card}} \leq S, \quad R_j \geq \frac{R_2^{j-1}}{S_{\text{card}}^{j-2}} \quad (R_2 > 0, j \geq 3).$$

Thus section sizes alone can give an explicit improvement over the cubic profile denominator; no search over the geometry of a section is needed.

Proof. Every nonempty coordinate product coset in $\mathbb{F}_p^k$ has cardinality p^a for some $0 \leq a \leq k$. Since $|A_r \triangle C| \geq ||A_r| - |C||$, we have $b_p(m_r) \leq \rho_r$. Both quantities lie in $[0, \tau(1/100)]$, where $x(1-x)$ is increasing. Equation (II.43.9) therefore gives the asserted bound for R_1 after summation. Insert it into (II.44.2) for the higher moments. $\square$

The two criteria can in fact be combined into a direct description of the original set and label, without keeping the auxiliary weights.

Theorem II.44.4 (Direct geometric classification of simultaneous equality). *Suppose $R_2 > 0$ and $j \geq 3$. Equality*

$$R_j = \frac{R_2^{j-1}}{S^{j-2}}$$

holds if and only if there are subgroups $H_i \leq G_i$, a subgroup $R' \leq R$, a point $r_0 \in R$, and representatives $c_r \in G$ for $r \in r_0 + R'$, with the following properties.

1. *The nonempty cross-sections are exactly those with $r \in r_0 + R'$, and each is $B_r = c_r + H$, where $H = \prod_i H_i$.*
2. *There is a separately additive map $M_0 : H \rightarrow K$ and a map $T : R' \times H_1 \times \cdots \times H_k \rightarrow K$ additive in all $k+1$ variables such that, for $x \in H$,*

$$\phi(c_r + x, r) = M_0(x) + T(r - r_0, x) + \sum_{U \subsetneq \{1, \dots, k\}} f_{r,U}(x_U), \quad (\text{II.44.5})$$

where all the lower-coordinate functions $f_{r,U}$ are arbitrary.

The representatives c_r may vary arbitrarily with r . Writing $m = |H|$ and $s = |R'|$, every such example satisfies

$$R_t = m^{2t+1} s^{2t-1} \quad (t \geq 1). \quad (\text{II.44.6})$$

Thus equality at one moment $j \geq 3$ implies equality at every higher moment.

Proof. By Corollary II.44.2, each active cross-section is $c_r + H_r$, where H_r is a coordinate product subgroup, and its full interaction is a separately additive map $M_r : H_r \rightarrow K$. Let $m_r = |H_r|$. A cube with sidelength h belongs to this section precisely when $h \in H_r$, and then it has m_r possible bases and the fixed label $M_r(h)$. Hence

$$a_h(r, v) = m_r \mathbf{1}_{h \in H_r} \mathbf{1}_{v = M_r(h)}.$$

At $h = 0$, the total mass of a_0 is $\sum_r m_r$. At any active h , it is $\sum_{r: h \in H_r} m_r$. The common-mass condition of Theorem II.44.1 forces these two sums to coincide. All summands are positive, so an element lying in any H_r lies in every H_r . The subgroups H_r are therefore all the same coordinate product subgroup H .

Let I be the set of active indices. Since $M_r(0) = 0$, the support of a_0 is $I \times \{0\}$. Its being a coset forces $I = r_0 + R'$ for a subgroup $R' \leq R$. For each $h \in H$, the support of a_h is the graph

$$\{(r, M_r(h)) : r \in r_0 + R'\}.$$

A graph over a subgroup coset is a subgroup coset in $R \times K$ exactly when its function is affine additive on that coset: translate the graph by its point at r_0 , then use closure under addition and the uniqueness of the graph value. Consequently

$$T(u, h) = M_{r_0+u}(h) - M_{r_0}(h)$$

is additive in $u \in R'$. For fixed u , it is a difference of two separately additive functions of h , so it is separately additive in the k remaining variables as well. Taking $M_0 = M_{r_0}$ and using Theorem II.43.5 gives (II.44.5).

Conversely, this description gives $a_h = m1_{\Gamma_h}$ for $h \in H$ and $a_h = 0$ otherwise, where each Γ_h is a coset graph of size s and every active mass is ms . Both equality criteria hold. More explicitly, there are m choices of h , s^{2t-1} index tuples with equal half sums, and m^{2t} choices of cube bases. The affine-additive dependence of $M_r(h)$ on r makes the label relation automatic. This proves (II.44.6). $\square$

[write] **6 October 2026; sources 15–20.** Source 15's Theorem II.54.2 is the same normal form, proved independently.

This classification allows very nonrectangular sets B as subsets of $G \times R$: the section translations may vary freely. If some H_i is trivial then M_0 and T vanish, and every labeling of each section is allowed. Such degenerate configurations are a necessary part of the exact arrangement count, rather than exceptions that may be discarded.

For $j = 8$, a known geometric defect supplies the explicit additional factor $(1 - \Xi/S)^{-6}$ relative to the cubic-profile moment bound. This is a local strengthening of Lemma 14.7's moment step. It does not by itself change the final quantitative bound for Szemerédi's theorem, because the later extraction and inductive losses must also be tracked.

II.44.1 A quantitative form of spectral flatness

Proposition II.44.5 (Near equality concentrates the active Fourier mass). *Let $R_2 > 0$, $j \geq 3$, and suppose*

$$R_j \leq (1 + \eta) \frac{R_2^{j-1}}{R_1^{j-2}}, \quad \eta \geq 0.$$

With X , ν , and μ as in the preceding proof,

$$\mathbb{E}_\nu \left(\frac{X}{\mu} - 1 \right)^2 = \frac{R_3 R_1}{R_2^2} - 1 \leq (1 + \eta)^{1/(j-2)} - 1. \quad (\text{II.44.7})$$

Thus, for every $s > 0$,

$$\sum_{|X/\mu-1| \geq s} \nu(h, \chi) \leq \frac{(1 + \eta)^{1/(j-2)} - 1}{s^2}. \quad (\text{II.44.8})$$

Proof. The identity in (II.44.7) follows by expanding the square. Interpolating the moments of the nonnegative spectrum gives

$$R_j \geq \frac{R_3^{j-2}}{R_2^{j-3}}.$$

For $j = 3$ this is equality. For $j > 3$, it follows directly from Hölder interpolation between the powers 2, 3, and j . Combining this inequality with the hypothesis yields

$$\left(\frac{R_3 R_1}{R_2^2} \right)^{j-2} \leq 1 + \eta.$$

Chebyshev's inequality proves (II.44.8). $\square$

[write] **6 October 2026; sources 15–20.** Source 15’s Theorem II.53.1 bounds the same variance by $\eta/(j-2)$, which this proposition improves; its Corollary II.53.2 and Proposition II.53.3 add the size of the near-flat band and the side-vector profile.

This conclusion concerns the Fourier mass weighted by its contribution to R_1 . It does not assert that the labeling is close to a separately additive function, or that every active weight a_h is close to a coset indicator. Those stronger inverse statements require additional work.

II.45 Source 15’s chapter: collision-sensitive amplification at Lemmas 14.6–14.8 [write]

Added in the fourth write, 6 October 2026.

The next eleven sections are source 15, a manuscript on Gowers’s Section 14: Corollary 14.6 turns the product property into respected pairs of parallel cubes, Lemma 14.7 amplifies them into respected eight-arrangements by moment interpolation, and Lemma 14.8 combines the two. Source 15 keeps the collision count $C = R_1$ that the second-moment estimate discards, and keeps it through both steps before combining them. *Logically the chapter follows source 04’s Section II.10* (arrangement moments with the cross-section profile) and sits beside source 14’s Sections II.43–II.44; it is printed here, at the end of Part II, because no number printed earlier may move. Its pin precedes the placement of both, so it is independent of them. Its introduction is Section F.20; its computations, Lean plan, questions, conclusion and appendices are in Part V (Sections V.59–V.63). Its notation is in the table at the head of this Part.

What is proved, in one paragraph. The interpolation $R_2^{d-1} \leq C^{d-2} R_d$ with its equality case, the cap $C \leq \sum_x b_x^3$ and the density-sensitive Lemma 14.7 are source 04’s (A5)–(A11) and source 14’s Theorem II.44.1, proved again (Theorem II.47.3, Lemma II.48.1, Corollary II.48.2). New: the product-property step with the collision count retained, $a_2 \geq \gamma^{8 \cdot 2^k} \chi^4$ (Theorem II.49.5), and the coupled bound $a_d \geq \gamma^{8(d-1)2^k} \chi^{3d-2} \geq \beta^{(3d-2)4^k} \gamma^{(6d-4)k4^k+8(d-1)2^k}$ (Theorem II.49.6); for eight-arrangements $\beta^{22 \cdot 4^k} \gamma^{44k4^k+56 \cdot 2^k}$, against $\beta^{28 \cdot 4^k} \gamma^{84k4^k}$ in the formal lemma_14_8 and $\beta^{28 \cdot 4^k-6} \gamma^{84k4^k}$ in source 04’s (A12), smaller in both exponents for every $k \geq 1$ (note there). Structure: the exact Cartesian classification with labels (source 14’s Theorem II.43.5), linear-error stability with 18ε (a weaker second route to source 14’s $\tau(\varepsilon)$), spectral stability (weaker than source 14’s Proposition II.44.5), the band and side-vector profile (new), the joint normal form (source 14’s Theorem II.44.4), a dichotomy and vertex weights (new). The source claims no improvement of a final bound and no priority, and does not claim the unsupported printed exponent $k4^{k+1}$.

Relation to the formal statements (checked at HEAD 0b4890a1c). corollary_14_6, lemma_14_7 and lemma_14_8 (Sections14_15.lean, lines 266, 279, 292) have exact companions (Proofs14Arrangements.lean, line 838; Proofs14HigherArrangements.lean, line 712; Proofs14ProductToArrangements.lean, line 18). The configuration input is formal (lemma_14_4_holds, corollary_14_5_holds, Proofs14Configurations.lean, lines 698, 942), as the source says; the unweakened Corollary 14.6 exponent is the private cor146_direct_lower (line 759), which the source calls documented (note in Section F.20). Nothing in this chapter is formalized; the source’s plan is in Part V and in 15-collision-amplification-integration.md.

II.46 Axis cubes and collision measures

Source 15, delivered Section 2.

Fix $k \geq 1$, finite abelian groups $G_1, \dots, G_k, T, V$, and write

$$X = \prod_{i=1}^k G_i, \quad M = |X|, \quad n = |T|, \quad q = |V|.$$

Let $B \subseteq X \times T$ and let $\phi : B \rightarrow V$ be a partial map. For $y, h \in X$ and $\omega \in \{0, 1\}^k$, the expression $y + \omega h$ means $(y_i + \omega_i h_i)_{i=1}^k$.

Definition II.46.1 (Admissible cube and its value). A triple $(y, h, x) \in X \times X \times T$ is admissible when $(y + \omega h, x) \in B$ for every $\omega \in \{0, 1\}^k$. For an admissible triple put

$$D_h \phi(y, x) = \sum_{\omega \in \{0, 1\}^k} (-1)^{k-|\omega|} \phi(y + \omega h, x). \quad (\text{II.46.1})$$

This is an alternating sum in the abelian group V .

The repository uses $(-1)^{|\omega|}$ instead. Its cube value differs by the fixed sign $(-1)^k$. Both equality and additive-respect conditions are unchanged. No extension of ϕ outside B affects any count.

Definition II.46.2 (Respected arrangement). For an integer $d \geq 1$, a respected d -arrangement consists of a common side vector $h \in X$ and $2d$ ordered admissible cubes (y_j, h, x_j) such that

$$\sum_{j=1}^d x_j = \sum_{j=d+1}^{2d} x_j, \quad (\text{II.46.2})$$

$$\sum_{j=1}^d D_h \phi(y_j, x_j) = \sum_{j=d+1}^{2d} D_h \phi(y_j, x_j). \quad (\text{II.46.3})$$

Let R_d count these arrangements, and let $R_d(h)$ count those with prescribed h .

In the unrestricted domain, ignoring (II.46.3), there are $M^{2d+1} n^{2d-1}$ arrangements: choose h , all $2d$ base points, and all but one of the vertical coordinates. This explains the normalization of a_d used in Section F.20. For $G_i = T = V = \mathbb{Z}/N\mathbb{Z}$, R_2 and R_8 are exactly the counts relevant to Lemmas 14.7–14.8, with ambient powers N^{5k+3} and N^{17k+15} .

Define a nonnegative integer-valued measure on $T \times V$ by

$$m_h(x, v) = \#\{y \in X : (y, h, x) \text{ is admissible and } D_h \phi(y, x) = v\}. \quad (\text{II.46.4})$$

The collision count and its two useful decompositions are

$$C = R_1 = \sum_{h, x, v} m_h(x, v)^2 = \sum_h C_h = \sum_x C_x, \quad (\text{II.46.5})$$

where $C_h = \sum_{x, v} m_h(x, v)^2$ and $C_x = \sum_{h, v} m_h(x, v)^2$. The subscripts h and x distinguish side-vector and vertical-fibre disintegrations.

For the fibre $B_x = \{y : (y, x) \in B\}$, put

$$b_x = |B_x|, \quad \beta_x = b_x/M, \quad \beta = \mathbb{E}_{x \in T} \beta_x, \quad \rho_3 = \mathbb{E}_{x \in T} \beta_x^3, \quad \Gamma = \sum_x b_x^3 = M^3 n \rho_3. \quad (\text{II.46.6})$$

When B is nonempty, $C > 0$: the zero side vector already supplies collisions. Empty domains are dealt with by the division-free statements.

II.47 Exact moment amplification

Source 15, delivered Section 3.

Let $\widehat{T}$ and $\widehat{V}$ denote the character groups, with characters valued in the unit circle. All Fourier transforms below are unnormalized.

Lemma II.47.1 (Moment-count identity). *For*

$$F_h(\alpha, \psi) = \sum_{x \in T, v \in V} m_h(x, v) \alpha(x) \psi(v), \quad \alpha \in \widehat{T}, \psi \in \widehat{V},$$

and every integer $d \geq 1$,

$$\sum_{h, \alpha, \psi} |F_h(\alpha, \psi)|^{2d} = nq R_d. \quad (\text{II.47.1})$$

Proof. Expand $F_h^d \overline{F_h}^d$. For an ordered selection $(x_1, v_1), \dots, (x_{2d}, v_{2d})$, its coefficient is $\prod_j m_h(x_j, v_j)$, and its character factor is

$$\alpha \left(\sum_{j \leq d} x_j - \sum_{j > d} x_j \right) \psi \left(\sum_{j \leq d} v_j - \sum_{j > d} v_j \right).$$

The sum over α, ψ is nq precisely when both arguments are zero, and is otherwise zero, by finite character orthogonality. The multiplicity $\prod_j m_h(x_j, v_j)$ chooses exactly the corresponding admissible base points. Summing over h proves the identity. This also proves the identity for general finite abelian value groups, without a prime-modulus assumption. $\square$

Lemma II.47.2 (Finite moment interpolation). *For finitely many nonnegative real numbers t_i and an integer $d \geq 2$,*

$$\left(\sum_i t_i^2 \right)^{d-1} \leq \left(\sum_i t_i \right)^{d-2} \sum_i t_i^d. \quad (\text{II.47.2})$$

If $d > 2$ and $\sum_i t_i > 0$, equality holds exactly when all positive t_i are equal.

Proof. If $S_1 = \sum_i t_i$ is zero, all terms vanish. Otherwise give index i probability $p_i = t_i/S_1$. Jensen's inequality for the convex function $u \mapsto u^{d-1}$ gives

$$\sum_i p_i t_i^{d-1} \geq \left(\sum_i p_i t_i \right)^{d-1}.$$

Multiplying by S_1^{d-1} yields (II.47.2). For $d > 2$, strict convexity says that equality holds precisely when t_i is constant on the support of p . For $d = 2$, the inequality is an identity. $\square$

Theorem II.47.3 (Collision-sensitive amplification). *For every integer $d \geq 2$,*

$$\boxed{R_2^{d-1} \leq C^{d-2} R_d.} \quad (\text{II.47.3})$$

If $C > 0$, this is equivalently

$$\boxed{a_d \geq \frac{a_2^{d-1}}{\chi^{d-2}.} \quad (\text{II.47.4})$$

For $d > 2$, equality holds exactly when the nonzero numbers $|F_h(\alpha, \psi)|^2$, across all three indices, have a common value.

Proof. Apply Lemma II.47.2 to $t_{h,\alpha,\psi} = |F_h(\alpha, \psi)|^2$. By Lemma II.47.1, the three sums in that inequality are nqR_2 , nqC , and nqR_d . The common factor $(nq)^{d-1}$ cancels. The powers of M and n also cancel in the stated normalization. The equality assertion is exactly the equality assertion of the interpolation lemma. If $C = 0$, every m_h vanishes, so all counts vanish. $\square$

[write] **6 October 2026; sources 15–20.** *Credit.* This is source 04’s interpolation (A5) (Theorem II.10.1) and source 14’s Theorem II.44.1, with the same equality criterion (source 14 proves it through the same coset-weight description, this source’s Lemma II.54.1); proved independently here, in a division-free form that also allows $C = 0$ and $d = 2$. *Formal status* (HEAD 0b4890a1c): Lemma II.47.1 for general d is the private `cor147_moment_eq_count` (Proofs14HigherArrangements.lean, line 484); the interpolation exists only at $d = 8$, in the private `cor147_interpolation` and `cor147_moment_interpolation` (lines 653 and 701, exponents 2, 4, 16), inside the proof of `lemma_14_7_holds` (line 712). The theorem for general d is not stated in Lean, and the “interpolation infrastructure” of Section V.59 is $d = 8$ only.

[write] **6 October 2026; sources 28–33.** For Gowers’s Lemma 12.3 source 33 adds a linear branch: θN^8 respected parallelogram pairs give at least $\max\{\theta^7, 4\theta - 3\}N^{32}$ respected eight-arrangements (Corollary II.137.6; the branch is active for $\theta > 0.8045\dots$), and a sharp repair theorem on full products (Section II.136).

Remark II.47.4 (Restriction of the side-vector family). The same proof works with h restricted to any fixed subset of X . It also works with a fixed nonnegative weight on h , provided the same weight is used in every moment and arrangement count. Excluding zero side coordinates alone therefore does not destroy the moment identity. Excluding arbitrary coincidences between different constituent cubes is another matter: such exclusions need not factor through m_h .

II.48 The cubic fibre bound and sharp density improvements

Source 15, delivered Section 4.

For $A \subseteq X$, define

$$Q(A) = \#\{(h, y, z) \in X^3 : y + \omega h, z + \omega h \in A \text{ for every } \omega \in \{0, 1\}^k\}. \quad (\text{II.48.1})$$

No label condition appears in Q .

Lemma II.48.1 (Cubic domination). *For every $A \subseteq X$ and every labelled domain B ,*

$$Q(A) \leq |A|^3, \quad C_x \leq Q(B_x) \leq b_x^3, \quad C \leq \Gamma. \quad (\text{II.48.2})$$

Consequently,

$$0 < \chi \leq \rho_3 \leq \beta \leq 1 \quad (\text{II.48.3})$$

for nonempty B .

Proof. The map

$$(h, y, z) \mapsto (y, y + h, z)$$

injects the configurations counted by $Q(A)$ into A^3 . Its inverse on its image recovers h as the second entry minus the first. Forgetting equality of the mixed values gives $C_x \leq Q(B_x)$. Sum over x . Finally, $0 \leq \beta_x \leq 1$ implies $\beta_x^3 \leq \beta_x$. $\square$

Corollary II.48.2 (Density-sensitive version of Lemma 14.7). *In the cyclic setting $G_i = T = V = \mathbb{Z}/N\mathbb{Z}$, suppose $R_2 \geq \theta N^{5k+3}$ with $\theta > 0$. Then*

$$R_8 \geq \theta^7 \chi^{-6} N^{17k+15} \geq \theta^7 \rho_3^{-6} N^{17k+15} \geq \theta^7 \beta^{-6} N^{17k+15}. \quad (\text{II.48.4})$$

If, in addition, $\beta_x \leq \kappa \beta$ for every x , where $\kappa \geq 1$, then

$$R_8 \geq \kappa^{-12} \theta^7 \beta^{-18} N^{17k+15}. \quad (\text{II.48.5})$$

Proof. Use Theorem II.47.3 with $d = 8$ and Lemma II.48.1. Under the fibre cap,

$$\rho_3 = \mathbb{E}_x \beta_x^3 \leq (\kappa \beta)^2 \mathbb{E}_x \beta_x = \kappa^2 \beta^3.$$

Substitute this upper bound into (II.48.4). □

[write] **6 October 2026; sources 15–20.** *Credit.* Lemma II.48.1 and this corollary are source 04's (A6)–(A11) (Theorem II.10.1): source 04's profile σ is this ρ_3 and its K is this κ ; the cubic cap $C \leq \sum_x b_x^3$ is source 04's (A9) and source 14's chain (II.43.1). Example II.48.3 below has source 04's (A13) and (A14) as its two faces $H = X$ and $J = T$.

For arbitrary d , the same argument gives

$$a_d \geq a_2^{d-1} \rho_3^{-(d-2)} \geq a_2^{d-1} \beta^{-(d-2)}, \quad (\text{II.48.6})$$

and the fibre cap gives the factor $\kappa^{-2(d-2)} \beta^{-3(d-2)}$. The density of a fibre matters cubically, not merely through its contribution to total mass.

Example II.48.3 (A sharp two-parameter family). Let $H_i \leq G_i$, put $H = \prod_i H_i$, let $J \leq T$, and take

$$B = (c + H) \times (t + J), \quad \phi = 0.$$

Writing $a = |H|/M$ and $b = |J|/n$, direct counting gives

$$\beta = ab, \quad \rho_3 = \chi = a^3 b, \quad a_d = a^{2d+1} b^{2d-1}. \quad (\text{II.48.7})$$

Indeed, an admissible side lies in H , every base lies in $c + H$, and there are $|J|^{2d-1}$ vertical additive tuples in $t + J$. Thus $R_d = |H|^{2d+1} |J|^{2d-1}$, and (II.47.4) is equality for every d . When $H = X$, the factor $\beta^{-(d-2)}$ in (II.48.6) is sharp. When $J = T$, all fibres have the same density, and the balanced factor $\beta^{-3(d-2)}$ is sharp.

Example II.48.4 (Why balance cannot be silently assumed). Take $B = X \times J$, where $J < T$ is a subgroup, and $\phi = 0$. Then $\rho_3 = \chi = \beta$, $a_2 = \beta^3$, and $a_8 = \beta^{15}$. Replacing ρ_3 by β^3 without a fibre hypothesis would assert $a_8 \geq a_2^7 / \beta^{18} = \beta^3$, which is false for $0 < \beta < 1$. Jensen's inequality actually gives $\rho_3 \geq \beta^3$, the opposite of the upper bound needed in that argument.

II.49 Coupling the product property to the collision count

Source 15, delivered Section 5.

This section proves (F.20.2). The important point is that χ enters the product-property lower bound with a positive fourth power before it enters the interpolation denominator. It must be retained in both places until these exponents have been combined.

II.49.1 The generalized product property

Let the $k + 1$ coordinate groups of $X \times T$ be denoted by $K_1, \dots, K_{k+1}$, with $K_{k+1} = T$.

Definition II.49.1 (Product property). The map $\phi : B \rightarrow V$ has PP_γ if the following holds for every coordinate j , every integer $p \geq 0$, every family of p parallel j -coordinate lines, every common subset $E \subseteq K_j$ on which all these lines lie in B , and every function $w : E \rightarrow [0, \infty)$. If $\psi_1, \dots, \psi_p : E \rightarrow V$ are the corresponding restrictions of ϕ , then

$$\sum_{\substack{u_1+u_2=u_3+u_4 \text{ in } E \\ \psi_i(u_1)+\psi_i(u_2)=\psi_i(u_3)+\psi_i(u_4) \ (1 \leq i \leq p)}} \prod_{s=1}^4 w(u_s) \geq \frac{\gamma^{8p}}{|K_j|} \left(\sum_{u \in E} w(u) \right)^4. \quad (\text{II.49.1})$$

Repeated lines are allowed. We assume $0 < \gamma \leq 1$.

For equal cyclic coordinate and value groups this is the repository's `HasProductProperty`, with its normalization unchanged [79]. Restriction to a coordinate fibre preserves the property: every relevant family of lines is still a family of lines in the original domain.

Lemma II.49.2 (An automatic weak parameter). *Every V -valued partial map on a product of finite abelian groups has PP_γ for $\gamma = q^{-1/8}$. A constant map has PP_1 .*

Proof. For a chosen family of p restrictions, place mass $w(u)$ at $(u, \psi_1(u), \dots, \psi_p(u)) \in K_j \times V^p$. Call this nonnegative measure μ . The left side of (II.49.1) is $\sum_z (\mu * \mu)(z)^2$. Cauchy–Schwarz bounds it below by

$$\frac{(\sum_z (\mu * \mu)(z))^2}{|K_j|q^p} = \frac{(\sum_u w(u))^4}{|K_j|q^p}.$$

This is (II.49.1) for the stated γ . For constant restrictions the value conditions are automatic, so the same argument on K_j alone gives $\gamma = 1$. $\square$

II.49.2 The configuration input, with all exponents retained

For a set $A \subseteq G_1 \times \dots \times G_j$, a configuration is specified by (y, h, t) and has vertices

$$y + \varepsilon h + \delta t, \quad \varepsilon, \delta \in \{0, 1\}^j.$$

It is contained if all vertices lie in A . It is respected if each coordinate parallelogram face satisfies the additive value identity: with all other binary coordinates fixed, the values at positions 00 and 11 sum to the values at positions 10 and 01. Let $\mathcal{F}_j(A, \phi)$ be the number of contained respected configurations.

Lemma II.49.3 (Configuration lower bound). *Suppose A has density α in a product of j finite abelian groups of total size L , and the restricted map has PP_γ . Then*

$$\mathcal{F}_j(A, \phi) \geq \alpha^{4j} \gamma^{2j4j} L^3. \quad (\text{II.49.2})$$

The configuration count is at most the respected parallel cube-pair count on that domain.

Proof. We prove the lower bound by induction, including $j = 0$. A product of zero groups has one point; its density is zero or one, so $\mathcal{F}_0 = \alpha$ and the assertion holds. For the induction step, write the ambient product as $Z \times K_j$, where $|Z| = L_0$ and $|K_j| = s$. For each potential $(j-1)$ -configuration ζ in Z , let E_ζ be the set of $u \in K_j$ at which ζ is contained and respected in the fibre A_u . Apply the product property to the 4^{j-1} vertical restrictions through its vertices, with weight $\mathbf{1}_{E_\zeta}$. The resulting additive quadruples supply exactly the final-coordinate faces; membership in E_ζ supplies all other faces. An additive quadruple is uniquely parameterized by its base and two side increments. Hence

$$\mathcal{F}_j(A, \phi) \geq \frac{\gamma^{8 \cdot 4^{j-1}}}{s} \sum_{\zeta} |E_\zeta|^4.$$

There are L_0^3 potential ζ . Also

$$\sum_{\zeta} |E_\zeta| = \sum_{u \in K_j} \mathcal{F}_{j-1}(A_u, \phi_u) \geq L_0^3 s \gamma^{e_{j-1}} \alpha^{4^{j-1}},$$

where $e_{j-1} = 2(j-1)4^{j-1}$; the last step uses the induction hypothesis and Jensen in the fibre densities. A further Jensen application gives

$$\mathcal{F}_j \geq \alpha^{4^j} \gamma^{4e_{j-1} + 8 \cdot 4^{j-1}} L_0^3 s^3.$$

Since $4e_{j-1} + 8 \cdot 4^{j-1} = 2j4^j$, this proves (II.49.2).

For the final assertion, send (y, h, t) to the two cubes with bases y and $y + t$, both with side h . The map is injective. Their mixed values are equal: move the base from y to $y + t$ one coordinate at a time. At coordinate i , the mixed-value change contains the face difference $\Delta_{h_i} \Delta_{t_i} \phi$, which vanishes at every required vertex by respect of the corresponding face. Every vertex used belongs to the configuration. Thus the image is a respected cube pair. $\square$

Corollary II.49.4 (Collision lower bound). *If $B \subseteq X \times T$ has density β and ϕ has PP_γ , then*

$$\chi \geq \beta^{4^k} \gamma^{e_k}, \quad e_k = 2k4^k. \quad (\text{II.49.3})$$

More precisely, $\chi \geq \gamma^{e_k} \mathbb{E}_x \beta_x^{4^k}$.

Proof. Apply Lemma II.49.3 to every horizontal fibre, where the product property is inherited. Its respected cube-pair count is C_x . Sum and divide by $M^3 n$. Jensen gives the displayed weaker form in β . $\square$

The argument above recovers the configuration and cube-pair input of Gowers's Lemma 14.4 and Corollary 14.5 in the present group notation. It is included to make the subsequent coupled bound self-contained, not as a claimed new configuration theorem.

II.49.3 The collision-parametrized pair bound

Theorem II.49.5 (Retaining the collision parameter). *Under PP_γ , put $u_k = 8 \cdot 2^k$. For every $h \in X$,*

$$R_2(h) \geq \frac{\gamma^{u_k} C_h^4}{n M^4}. \quad (\text{II.49.4})$$

Consequently,

$$R_2 \geq \frac{\gamma^{u_k} C^4}{nM^7}, \quad \boxed{a_2 \geq \gamma^{u_k} \chi^4.} \quad (\text{II.49.5})$$

Only the product property in the vertical coordinate is needed here.

Proof. Fix h and a reference base $y \in X$. For each $x \in T$, define

$$w_y(x) = \begin{cases} m_h(x, D_h \phi(y, x)), & (y, h, x) \text{ admissible,} \\ 0, & \text{otherwise.} \end{cases}$$

This counts admissible cubes at level x with the same side and mixed value as the reference cube. Let E_y be the set of levels where the reference cube is admissible. Apply (II.49.1) to the 2^k vertical lines through $y + \omega h$, the common domain E_y , and the nonnegative weight w_y . It gives at least

$$\frac{\gamma^{u_k}}{n} \left(\sum_x w_y(x) \right)^4$$

weighted additive quadruples of levels whose values are additive at every vertex of the reference cube. Alternating over the vertices shows that the reference mixed values also satisfy the additive relation.

For each of the four levels, choose a cube counted by w_y . Its mixed value equals the corresponding reference mixed value, so the four chosen cubes form a respected 2-arrangement with side h . For this fixed reference base y , the resulting choices inject into the arrangements counted by $R_2(h)$: the levels and the four chosen bases recover the choices. Therefore

$$R_2(h) \geq \frac{\gamma^{u_k}}{n} \left(\sum_x w_y(x) \right)^4 \quad \text{for each } y.$$

Average over all M reference bases and apply Jensen. The exact identity

$$\sum_{y,x} w_y(x) = \sum_{x,v} m_h(x, v)^2 = C_h$$

then gives (II.49.4). Notice that the injection is asserted for a *fixed* reference base. Averaging, rather than claiming a global injection without that base, accounts for the factor M .

Finally sum over the M side vectors and use $\sum_h C_h^4 \geq C^4/M^3$. This proves (II.49.5) and its normalized form. $\square$

[write] **6 October 2026; sources 15–20.** *Formal status* (HEAD 0b4890a1c). The ingredients are proof-level facts of the development: the first moment $\sum \theta = C$ is the right side of the private `cor146_pair_first_moment` (`Proofs14Arrangements.lean`, line 183), and the private `cor146_average_fourth_moment` takes the first-moment lower bound as a hypothesis; the C -parametrized form proved here is not stated in Lean. Not formalized; new as a stated theorem.

Theorem II.49.6 (Coupled amplification). *Suppose $\phi : B \rightarrow V$ has PP_γ and B has density $\beta > 0$. For every integer $d \geq 2$,*

$$a_d \geq \gamma^{u_k(d-1)} \chi^{3d-2}, \quad (\text{II.49.6})$$

$$a_d \geq \beta^{(3d-2)4^k} \gamma^{(6d-4)k4^k + 8(d-1)2^k}. \quad (\text{II.49.7})$$

For $d = 8$, this becomes (F.20.3).

Proof. Since $C > 0$, combine Theorems II.47.3 and II.49.5 before estimating χ :

$$a_d \geq \frac{a_2^{d-1}}{\chi^{d-2}} \geq \frac{(\gamma^{u_k} \chi^4)^{d-1}}{\chi^{d-2}} = \gamma^{u_k(d-1)} \chi^{3d-2}.$$

The remaining power of χ is positive. We may now apply the *lower* bound (II.49.3). The exponent of γ is

$$(3d-2)e_k + (d-1)u_k = (6d-4)k4^k + 8(d-1)2^k,$$

which proves (II.49.7). □

[write] 6 October 2026; sources 15–20. *New in this report, and stronger than the formal bound* (exponent arithmetic checked at the write for $1 \leq k \leq 40$, $2 \leq d \leq 11$). At $d = 8$ the bound (F.20.3) is $\beta^{22 \cdot 4^k} \gamma^{44k4^k + 56 \cdot 2^k}$, against $\beta^{28 \cdot 4^k} \gamma^{84k4^k}$ in `lemma_14_8` (Sections14_15.lean, line 292: $\beta^{7 \cdot 4^{k+1}} \gamma^{21k4^{k+1}}$; exact companion `lemma_14_8_holds`, Proofs14ProductToArrangements.lean, line 18) and $\beta^{28 \cdot 4^k - 6} \gamma^{84k4^k}$ in source 04's (A12). Both exponents are smaller for every $k \geq 1$: $22 \cdot 4^k < 28 \cdot 4^k - 6$, and $44k4^k + 56 \cdot 2^k \leq 84k4^k$ because $56 \cdot 2^k \leq 40k4^k$. Against the unweakened corrected exponent $(7(2k4^{k+1} + 8 \cdot 2^k))$, proof-level in `cor146_direct_lower`, note in Section F.20) the gain is exactly the factor $(\beta^{4^k} \gamma^{2k4^k})^{-6}$, as Section II.50 says. No other source of this report couples the two steps. Not formalized.

Remark II.49.7 (A direction-of-inequality trap). It would be invalid to replace the denominator χ^{d-2} by a *lower* bound while independently replacing a_2 by a lower bound. The argument succeeds because it first substitutes the functional inequality $a_2 \geq \gamma^{u_k} \chi^4$, leaving a positive power of the same χ . This is why the gain is larger than simply replacing the ambient second-moment cap by a density cap.

A finer, immediately available version of (II.49.7) replaces $\beta^{(3d-2)4^k}$ by $(\mathbb{E}_x \beta_x^{4^k})^{3d-2}$. The theorem does not claim that its exponents are optimal. The sharpness examples in Section II.48 establish sharpness of the collision interpolation and the cubic-density substitutions, not of every product-property estimate used here.

II.50 Quantitative comparison with the supplied statements

Source 15, delivered Section 6.

For this section specialize to $G_i = T = V = \mathbb{Z}/N\mathbb{Z}$. Theorem II.49.5 and Corollary II.49.4 give the unweakened pair estimate

$$a_2 \geq \beta^{4^{k+1}} \gamma^{2k4^{k+1} + 8 \cdot 2^k}. \quad (\text{II.50.1})$$

This derives the exponent E_k recorded in the source correction rather than assuming the printed smaller exponent. For $k \geq 1$,

$$E_k \leq 3k4^{k+1} \iff 8 \cdot 2^k \leq k4^{k+1},$$

which holds, with equality at $k = 1$.

Applying the ambient version of Lemma 14.7 to the repository's safe pair bound yields its supported Lemma 14.8:

$$a_8 \geq \beta^{28 \cdot 4^k} \gamma^{84k4^k}. \quad (\text{II.50.2})$$

A density-only substitution in Corollary II.48.2 improves this to

$$a_8 \geq \beta^{28 \cdot 4^k - 6} \gamma^{84k4^k}. \quad (\text{II.50.3})$$

With uniform fibre density, the subtraction from the β exponent is 18 instead of 6. These refinements need no product-property argument beyond its already supplied lower bound on a_2 .

The coupled result is stronger:

$$a_8 \geq \beta^{22 \cdot 4^k} \gamma^{44k4^k + 56 \cdot 2^k}. \quad (\text{II.50.4})$$

Because $0 < \beta, \gamma \leq 1$, decreasing either exponent increases the lower bound. The comparisons in Table II.50.1 are therefore pointwise improvements, not merely asymptotic ones.

Table II.50.1: Exponents (b, g) in a proved lower bound $a_8 \geq \beta^b \gamma^g$. The middle column uses the corrected safe statement, not the unsupported printed γ exponent.

k	Repository's corrected Lemma 14.8	Coupled theorem
1	(112, 336)	(88, 288)
2	(448, 2688)	(352, 1632)
3	(1792, 16128)	(1408, 8896)

There are two distinct sources of savings, which should not be conflated. Retaining the already documented E_k instead of its safe upper estimate changes the old γ exponent from $84k4^k$ to $56k4^k + 56 \cdot 2^k$. The additional collision cancellation changes

$$(28 \cdot 4^k, 56k4^k + 56 \cdot 2^k) \quad \text{to} \quad (22 \cdot 4^k, 44k4^k + 56 \cdot 2^k).$$

Thus the second, genuinely coupled saving is six powers of the collision lower bound $\beta^{4^k} \gamma^{2k4^k}$.

No source claim of an “exact companion” is used as a substitute for checking its type. The supplied definition and proof of Lemma 14.8 actually use the repaired exponent [79, 82]. The present comparison is with those types. The general-group version of the result also does not follow merely from changing a type name in the Lean code: its proof here uses the generalized product property and character orthogonality stated above.

II.51 Equality in the collision cap

Source 15, delivered Section 7.

We now determine the geometric meaning of saturating the cubic estimate. For a coordinate group product X , write $\pi_i : X \rightarrow X$ for projection onto coordinate i , with all other coordinates set to zero.

Theorem II.51.1 (Exact Cartesian classification). *For a nonempty set $A \subseteq X$, the following are equivalent:*

- (i) $Q(A) = |A|^3$;
- (ii) $A = c + H$ for a subgroup $H \leq X$ invariant under every π_i ;
- (iii) $A = c + (H_1 \times \cdots \times H_k)$ for subgroups $H_i \leq G_i$.

Proof. Equality in the injection of Lemma II.48.1 makes that injection surjective. Hence for every $r, t, s \in A$, the two cubes with common side $t - r$ and bases r, s lie in A . In particular their all-one corner gives $s + t - r \in A$. Choose $c \in A$ and let $H = A - c$. This set contains zero and is closed

under the ternary operation $u + v - w$; it is consequently a subgroup, using that operation with one or two inputs zero.

The single-coordinate corners also give $r + \pi_i(t - r) \in A$. Taking $r = c$ shows $\pi_i H \subseteq H$. If H_i is the i th coordinate projection of H , then every element of H belongs to $\prod_i H_i$, and every coordinate-supported element from H_i belongs to H . Adding them proves $H = \prod_i H_i$.

Conversely, in a Cartesian coset every triple $(r, t, s) \in A^3$ gives admissible cubes with side $t - r$: every coordinate of every vertex remains in its corresponding coordinate coset. The injection is therefore surjective, proving $Q(A) = |A|^3$. $\square$

The label condition has an equally explicit description. An arbitrary function that omits a coordinate has zero full mixed derivative, so such functions must be allowed in any correct classification.

Theorem II.51.2 (Labels at exact collision equality). *Let $A \neq \emptyset$ and $\phi : A \rightarrow V$. The respected parallel cube-pair count on A equals $|A|^3$ if and only if A is a Cartesian coset and, after translating it to $H = \prod_i H_i$, the map has the form*

$$\phi(x_1, \dots, x_k) = L(x_1, \dots, x_k) + \sum_{i=1}^k \psi_i(x_1, \dots, \hat{x}_i, \dots, x_k), \quad (\text{II.51.1})$$

where $L : H \rightarrow V$ is additive separately in every coordinate and the ψ_i are arbitrary functions on the indicated smaller products. For $k = 1$, the final sum is simply a constant.

Proof. Equality forces equality in both inequalities $C_A \leq Q(A) \leq |A|^3$, so Theorem II.51.1 applies. For each $h \in H$, every base is admissible. Equality of all collisions therefore says that $D_h \phi(y)$ is independent of $y \in H$.

Define the anchored mixed part by

$$L(x) = D_x \phi(0) = \sum_{S \subseteq [k]} (-1)^{k-|S|} \phi(x_S), \quad (\text{II.51.2})$$

where x_S agrees with x on S and is zero elsewhere. For a coordinate translation T_a and difference $\Delta_a = T_a - I$, the identity

$$\Delta_{a+b} = \Delta_a + T_a \Delta_b$$

holds. Apply it to one of the coordinate increments in (II.51.2), keeping the other differences fixed. The translated full mixed difference equals its unshifted value because of the base-point independence just proved. Thus L is additive in that coordinate. The same argument applies to every coordinate.

Solving (II.51.2) for the term with $S = [k]$ gives

$$\phi(x) - L(x) = \sum_{S \subsetneq [k]} (-1)^{k-|S|+1} \phi(x_S).$$

Every summand omits at least one coordinate. Assign each proper S to one chosen missing coordinate and collect the assigned terms to obtain the functions ψ_i in (II.51.1).

Conversely, the full mixed derivative of a function omitting a coordinate is zero. Separate additivity gives $D_h L(y) = L(h)$, independently of y . Hence every admissible pair with the same side has equal mixed values, and the count is $Q(A) = |A|^3$. $\square$

[write] **6 October 2026; sources 15–20.** *Credit.* Theorems II.51.1 and II.51.2 are source 14’s Theorem II.43.5 (equality with labels), which adds that the top interaction is unique; proved independently. At $k = 1$ the set statement is source 11’s Lemma II.23.3.

No division by factorials or scalar-field hypothesis enters this proof. It applies to arbitrary finite abelian value groups, including groups with small torsion. It describes a multiadditive top component modulo coordinate-omitting functions, not necessarily an ordinary polynomial over a field. Globally, $C = \Gamma$ holds precisely when every nonempty fibre satisfies this description, with possibly different subgroups and functions in different fibres.

II.52 Dimension-free Cartesian stability

Source 15, delivered Section 8.

The exact theorem has a robust counterpart with a linear error term. The proof first finds an almost-period subgroup using ordinary additive energy, then uses the axis-cube condition to force that subgroup to split by coordinates. The latter step is what ordinary energy alone would not give.

Lemma II.52.1 (A small-difference-set fact). *If a nonempty finite subset P of an abelian group satisfies $|P - P| < \frac{3}{2}|P|$, then $P - P$ is a subgroup.*

Proof. Put $D = P - P$. For $t = u - v \in D$, the sets $P - u$ and $P - v$ are both subsets of D , so their intersection has size at least $2|P| - |D| > |P|/2$. Thus

$$r_P(t) := |\{x \in P : x + t \in P\}| > |P|/2.$$

For any $s, t \in D$, the two subsets $\{x \in P : x + s \in P\}$ and $\{x \in P : x - t \in P\}$ each contain more than half of P and therefore intersect. For a point in the intersection, $(x + s) - (x - t) = s + t \in D$. Since D also contains zero and is symmetric, it is a subgroup. $\square$

[write] **6 October 2026; sources 15–20.** *Credit.* The same classical fact is source 14’s Lemma II.43.1 (note there); both prove it.

Theorem II.52.2 (Linear-error Cartesian stability). *Let $A \neq \emptyset$ be a subset of $X = \prod_{i=1}^k G_i$, and put $a = |A|$. Suppose*

$$Q(A) \geq (1 - \varepsilon)a^3, \quad 0 \leq \varepsilon \leq \frac{1}{100}.$$

There is a Cartesian coset $\mathcal{C} = c + (H_1 \times \cdots \times H_k)$ such that

$$|A \cap \mathcal{C}| \geq (1 - 8\varepsilon)a, \quad |\mathcal{C}| \leq (1 + 2\varepsilon)a, \quad |A \triangle \mathcal{C}| \leq 18\varepsilon a. \quad (\text{II.52.1})$$

All constants are independent of k and of the coordinate-group orders. The same conclusion follows if the respected cube-pair count for any labelling of A is at least $(1 - \varepsilon)a^3$.

Proof. Step 1: almost maximal additive energy. Let $r(t) = |\{x \in A : x + t \in A\}|$. We have $\sum_t r(t) = a^2$. Every cube pair counted by $Q(A)$ has its four extreme endpoints in A , so

$$Q(A) \leq E(A) := \sum_t r(t)^2 \leq a^3.$$

Consequently

$$\sum_t r(t)(a - r(t)) = a^3 - E(A) \leq \varepsilon a^3. \quad (\text{II.52.2})$$

Step 2: a subgroup of popular differences. Set $P = \{t : r(t) > 7a/8\}$. This set contains zero and is symmetric. For $t \notin P$, $r(t)(a - r(t)) \geq (a/8)r(t)$, so

$$\sum_{t \in P} r(t) \geq (1 - 8\varepsilon)a^2, \quad |P| \geq (1 - 8\varepsilon)a. \quad (\text{II.52.3})$$

The translation distance $a - r(t) = \frac{1}{2}|A \triangle (A - t)|$ obeys the triangle inequality. Hence for $s, t \in P$, $r(s - t) > 3a/4$. It follows that

$$|P - P| < \frac{4}{3}a < \frac{3}{2}(1 - 8\varepsilon)a \leq \frac{3}{2}|P|,$$

where the strict middle inequality uses $\varepsilon \leq 1/100$. By Lemma II.52.1, $H = P - P$ is a subgroup. Write $m = |H|$; then $m < 4a/3$ and $P \subseteq H$.

Step 3: one coset captures nearly all of A. The identity

$$\sum_{h \in H} r(h) = \sum_{\mathcal{D} \in X/H} |A \cap \mathcal{D}|^2$$

and (II.52.3) imply that one coset $\mathcal{C} = c + H$ has intersection size $b = |A \cap \mathcal{C}| \geq (1 - 8\varepsilon)a$. For every $h \in H$, two subsets of the same m -point coset give

$$r(h) \geq 2b - m > \left(\frac{2}{3} - 16\varepsilon\right)a \geq \frac{a}{2}.$$

Use this lower bound in (II.52.2), restricting the sum to H :

$$\begin{aligned} \varepsilon a^3 &\geq \frac{a}{2} \left(ma - \sum_{h \in H} r(h) \right) \\ &= \frac{a}{2} \left(ma - \sum_{\mathcal{D} \in X/H} |A \cap \mathcal{D}|^2 \right) \geq \frac{a^2}{2} (m - a). \end{aligned}$$

The final inequality uses $\sum_{\mathcal{D}} |A \cap \mathcal{D}|^2 \leq a^2$. It follows that $m \leq (1 + 2\varepsilon)a$ and

$$|A \triangle \mathcal{C}| = a + m - 2b \leq 18\varepsilon a. \quad (\text{II.52.4})$$

At this point H need not yet be a Cartesian subgroup.

Step 4: the axis condition forces a Cartesian subgroup. Fix a coordinate projection π_i and define $L_i(r, t) = r + \pi_i(t - r)$. The injection underlying $Q(A) \leq a^3$ shows that at least $(1 - \varepsilon)a^2$ ordered pairs $(r, t) \in A^2$ satisfy $L_i(r, t) \in A$: a pair failing this condition can support none of the at most a third entries in that injection. Thus at least $b^2 - \varepsilon a^2$ pairs from $(A \cap \mathcal{C})^2$ satisfy this condition.

For uniformly chosen $(r, t) \in \mathcal{C}^2$, the value $L_i(r, t)$ is uniform on the coset $c + K_i$, where

$$K_i = (I - \pi_i)H + \pi_i H = H + \pi_i H.$$

If $\pi_i H \not\subseteq H$, then $|K_i| \geq 2m$, and the number of pairs in $\mathcal{C}^2$ with $L_i(r, t) \in A$ is at most $am/2$. On the other hand,

$$b^2 - \varepsilon a^2 \geq ((1 - 8\varepsilon)^2 - \varepsilon)a^2 \geq (1 - 17\varepsilon)a^2 \geq \frac{83}{100}a^2,$$

whereas

$$\frac{am}{2} \leq \frac{1+2\varepsilon}{2}a^2 \leq \frac{51}{100}a^2.$$

This is a contradiction. Hence $\pi_i H \subseteq H$ for every i , and H is a product subgroup by the argument in Theorem II.51.1. This proves (II.52.1). The labelled version follows because its count is at most $Q(A)$. $\square$

[write] **6 October 2026; sources 15–20.** A weaker second route (checked at the write). Source 14's Theorem II.43.3 has the same hypothesis and gives a unique closest coordinate-product coset $\mathcal{C}$ with $|A \triangle \mathcal{C}| \leq \tau(\varepsilon)a$, $\tau(\varepsilon) = (1 - \sqrt{1 - 4\varepsilon})/2$. Since $\tau(\varepsilon)/\varepsilon$ increases on $[0, 1/100]$ and $\tau(1/100) = 0.010102\dots$, $\tau(\varepsilon) \leq 1.0102\varepsilon \leq 2\varepsilon$ there, so source 14's coset satisfies all three bounds of (II.52.1): $|A \cap \mathcal{C}| \geq (1 - \tau)a$, $|\mathcal{C}| \leq (1 + \tau)a$, $|A \triangle \mathcal{C}| \leq \tau a$. The theorem here is correct; its constants are not sharp. The method is the same (this source's Lemma II.52.1 is source 14's three-halves lemma, and the coordinate-mixing step is source 14's), and Example II.52.3 is source 14's Proposition II.43.4.

Example II.52.3 (The linear order is necessary). Let $k = 1$, take $G_1 = \mathbb{Z}/p\mathbb{Z}$ with $p \geq 101$ prime, and let $A = G_1 \setminus \{0\}$. Put $a = p - 1$. Then $r(0) = a$ and $r(t) = a - 1$ for every $t \neq 0$. Since $Q(A) = E(A)$ in dimension one,

$$Q(A) = a^2 + a(a - 1)^2 = a^3 - a^2 + a, \quad \varepsilon = \frac{1}{a} - \frac{1}{a^2}.$$

The closest coset of a subgroup is the whole group, at symmetric-difference distance 1. Its relative distance $1/a$ is asymptotic to ε . Thus no uniform $o(\varepsilon)$ replacement for the error term is possible. Taking the product with whole additional coordinate groups gives the same relative deficit and relative distance in higher dimensions. This does not establish optimality of the numerical constant 18.

[write] **6 October 2026; sources 15–20.** With source 14's theorem this example settles the leading constant asked for in the first question of Section V.60: at $\varepsilon = 1/a - 1/a^2$ one has $1 - 4\varepsilon = (1 - 2/a)^2$, so $\tau(\varepsilon)a = 1$, the distance here. The optimal $\delta(\varepsilon)$ is therefore $\tau(\varepsilon)$ along this sequence, with leading constant 1 (item II.31).

II.53 Stability of the amplification inequality

Source 15, delivered Section 9.

The preceding theorem describes near equality in the *geometric cap* $C \leq \Gamma$. Near equality in the *moment interpolation* has a different consequence: the spectrum has nearly constant magnitude on the part that carries its collision mass.

Assume $C > 0$ and put, with $i = (h, \alpha, \psi)$,

$$t_i = |F_h(\alpha, \psi)|^2, \quad p_i = \frac{t_i}{nqC}, \quad \lambda = \frac{R_2}{C}, \quad W_i = \frac{t_i}{\lambda}. \quad (\text{II.53.1})$$

Then (p_i) is a probability distribution and $\mathbb{E}_p W = 1$.

Theorem II.53.1 (Spectral plateau stability). For $d \geq 3$, define the nonnegative interpolation defect

$$\mathcal{D}_d = \frac{R_d C^{d-2}}{R_2^{d-1}} - 1.$$

Then

$$\mathbb{E}_p(W - 1)^2 \leq \frac{\mathcal{D}_d}{d - 2}. \quad (\text{II.53.2})$$

For every $\tau > 0$,

$$p\{i : |W_i - 1| \geq \tau\} \leq \frac{\mathcal{D}_d}{(d - 2)\tau^2}. \quad (\text{II.53.3})$$

Proof. The moment identity gives $\mathcal{D}_d = \mathbb{E}_p W^{d-1} - 1$. For $w \geq 0$ the exact polynomial identity

$$w^{d-1} - 1 - (d - 1)(w - 1) = (w - 1)^2 \sum_{j=0}^{d-3} (d - 2 - j)w^j \geq (d - 2)(w - 1)^2 \quad (\text{II.53.4})$$

holds. Taking expectations and using $\mathbb{E}_p W = 1$ proves (II.53.2). Markov's inequality applied to $(W - 1)^2$ proves (II.53.3). $\square$

[write] 6 October 2026; sources 15–20. *Credit.* Source 14's Proposition II.44.5 bounds the same variance (in its notation $\eta = \mathcal{D}_d$) by $(1 + \eta)^{1/(d-2)} - 1$, which is at most $\eta/(d - 2)$ by Bernoulli's inequality; so this theorem is a weaker second route (at $d = 8$, $\mathcal{D}_8 = 1$: 0.1225 against 1/6). Corollary II.53.2, Proposition II.53.3 and Example II.53.4 are new in this report.

For eight-arrangements, the coefficient in (II.53.2) is 1/6. This estimate concerns collision-weighted spectral mass, not the uniform proportion of all frequencies. Small or zero Fourier coefficients can occupy many indices without affecting it appreciably.

Corollary II.53.2 (Size of the near-flat band). *Define*

$$K_{\text{eff}} = \frac{nqC^2}{R_2}, \quad I_\tau = \{i : |W_i - 1| \leq \tau\}, \quad 0 < \tau < 1,$$

and let $\eta = \mathcal{D}_d/((d - 2)\tau^2)$. Then

$$\frac{\max(0, 1 - \eta)}{1 + \tau} K_{\text{eff}} \leq |I_\tau| \leq \frac{K_{\text{eff}}}{1 - \tau}. \quad (\text{II.53.5})$$

Proof. On I_τ , $\lambda(1 - \tau) \leq t_i \leq \lambda(1 + \tau)$. Its p -mass is at least $1 - \eta$, while its total t -mass is at most nqC . Divide these two bounds by the appropriate endpoints of the band. Since $nqC/\lambda = K_{\text{eff}}$, the conclusion follows. $\square$

II.53.1 The additional gain from the side-vector profile

Proposition II.53.3 (Disintegrated amplification). *For $d \geq 2$,*

$$R_d \geq \sum_{h:C_h>0} \frac{R_2(h)^{d-1}}{C_h^{d-2}} \geq \frac{R_2^{d-1}}{C^{d-2}}. \quad (\text{II.53.6})$$

For $C > 0$, put $w_h = C_h/C$, $\lambda_h = R_2(h)/C_h$ on occupied sides, and $\lambda = R_2/C$. The ratio of the middle expression in (II.53.6) to the final expression equals

$$\mathbb{E}_w \left(\frac{\lambda_h}{\lambda} \right)^{d-1}.$$

For $d \geq 3$, this ratio is at least

$$1 + (d - 2) \text{Var}_w \left(\frac{\lambda_h}{\lambda} \right). \quad (\text{II.53.7})$$

Proof. Apply the moment interpolation with h fixed and sum over h . Since $R_2(h) = 0$ whenever $C_h = 0$, no term is lost. The middle expression is $C\mathbb{E}_w\lambda_h^{d-1}$, and $\mathbb{E}_w\lambda_h = \lambda$. Jensen proves the second inequality and the ratio identity. Formula (II.53.4) proves the variance bound. $\square$

The full spectral distribution disintegrates over h with weights w_h . The law of total variance gives the useful identity

$$\mathrm{Var}_p(W) = \mathbb{E}_w \mathrm{Var}(W \mid h) + \mathrm{Var}_w(\lambda_h/\lambda).$$

Thus a near extremizer must have both approximately flat conditional spectra and approximately aligned scales across the collision-carrying side vectors.

Example II.53.4 (Spectral equality is not Cartesian equality). Let $X = (\mathbb{Z}/p\mathbb{Z})^2$, let T and V be trivial groups, and take the horizontal domain $A = \{(u, u) : u \in \mathbb{Z}/p\mathbb{Z}\}$ with the zero label. The only admissible side vector is $h = 0$, because a nonzero single-coordinate corner leaves the diagonal. Hence $R_d = p^{2d}$ and $C = p^2$. Moment amplification is an equality for every d , and the spectrum has one nonzero magnitude. But A is not a Cartesian coset for $p > 1$, and $Q(A)/|A|^3 = 1/p$. Near equality in the interpolation alone therefore cannot replace the hypothesis of Theorem II.52.2.

II.54 A complete normal form for joint equality

Source 15, delivered Section 10.

There is a useful stronger conclusion when both the moment inequality and the cubic collision cap are equalities. We first translate flat Fourier magnitudes into a statement about the nonnegative measures m_h .

Lemma II.54.1 (Positive flat spectra). *Let K be a finite abelian group and $\mu : K \rightarrow [0, \infty)$ a nonzero function of total mass s . Every Fourier coefficient has magnitude either 0 or s if and only if*

$$\mu = \frac{s}{|L|} \mathbf{1}_{z+L}$$

for some subgroup $L \leq K$ and some coset $z + L$.

Proof. Choose $z \in \mathrm{supp} \mu$ and let L be the subgroup generated by $\mathrm{supp} \mu - \mathrm{supp} \mu$. Equality in the triangle inequality $|\hat{\mu}(\xi)| \leq s$ holds exactly when ξ is constant on $\mathrm{supp} \mu$, equivalently when $\xi \in L^\perp$. For such characters the coefficient is $s\xi(z)$; for other characters, the hypothesis forces it to be zero. Fourier inversion and character orthogonality therefore give $(s/|L|)\mathbf{1}_{z+L}$. The converse follows by summing a character over that coset. $\square$

For a nonempty domain and $d \geq 3$, equality in moment amplification is therefore equivalent to the following description: every occupied m_h is a constant multiple of a subgroup-coset indicator in $T \times V$, and *all occupied measures have the same total mass*. The common mass condition is needed because the interpolation combines different h .

Theorem II.54.2 (Joint extremizers). *For a nonempty labelled domain $B \subseteq X \times T$, the following are equivalent:*

- (i) $C = \Gamma$ and $R_d C^{d-2} = R_2^{d-1}$ for some integer $d \geq 3$;

(ii) there exist a Cartesian subgroup $H = \prod_i H_i \leq X$, a subgroup $J \leq T$, a coset $S = x_0 + J$, and representatives $c_x \in X$ for $x \in S$ such that

$$B_x = c_x + H \quad (x \in S), \quad B_x = \emptyset \quad (x \notin S), \quad (\text{II.54.1})$$

and, for $u \in H$ and $x \in S$,

$$\phi(c_x + u, x) = L_0(u) + A(u, x - x_0) + \sum_{i=1}^k \psi_{i,x}(u_1, \dots, \hat{u}_i, \dots, u_k). \quad (\text{II.54.2})$$

Here $L_0 : H \rightarrow V$ is separately additive in the k horizontal variables, $A : H \times J \rightarrow V$ is separately additive in all $k + 1$ variables, and the functions $\psi_{i,x}$ are arbitrary.

For all these domains and every integer $r \geq 1$,

$$R_r = |H|^{2r+1} |J|^{2r-1}. \quad (\text{II.54.3})$$

In particular, equality then holds in moment amplification for every $r \geq 2$.

Proof. Assume (i). Since $k \geq 1$, the zero side has zero mixed value and $m_0(x, v) = b_x \mathbf{1}_{\{v=0\}}$. By Lemma II.54.1 and the moment equality, the nonempty vertical fibres form a coset $S = x_0 + J$, and all their sizes equal some positive integer b . The total mass of m_0 is $b|S|$, so every other occupied m_h has this same mass.

The equality $C = \Gamma$ holds fibre by fibre. By Theorems II.51.1 and II.51.2, each B_x is $c_x + H_x$ for a Cartesian subgroup H_x , and its mixed value $L_x(h)$ is independent of the base whenever $h \in H_x$. For any h that lies in at least one H_x ,

$$\sum_{x,v} m_h(x, v) = b |\{x \in S : h \in H_x\}|.$$

This is an occupied side, so its mass must equal $b|S|$. Thus such an h lies in every H_x , proving that all the subgroups H_x coincide with one Cartesian subgroup H .

For each $h \in H$, the measure m_h is now b times the indicator of

$$\{(x, L_x(h)) : x \in S\} \subseteq T \times V.$$

By Lemma II.54.1, this graph is a subgroup coset. Projection onto T is one-to-one on it, so its subgroup is the graph of a homomorphism from J to V . Consequently

$$L_{x_0+t}(h) = L_{x_0}(h) + A(h, t), \quad t \in J,$$

where $A(h, \cdot)$ is additive. Each L_x is separately additive in the horizontal variables by Theorem II.51.2; hence their difference $A(\cdot, t)$ is also separately additive. Put $L_0 = L_{x_0}$ and use the coordinate-omitting decomposition in each fibre to obtain (II.54.2).

Conversely, suppose (ii) holds. For $h \in H$, every base in every occupied fibre gives an admissible cube, and its mixed value is $L_0(h) + A(h, x - x_0)$; for $h \notin H$, no cube is admissible. Whenever the vertical coordinates satisfy the additive relation, these mixed values also satisfy it, by additivity of A in its last variable. Thus all vertical additive tuples are respected. Choosing the common side, the $2r$ bases, and the vertical tuple gives (II.54.3). Also $C = |H|^3 |J| = \Gamma$, and substituting these counts proves all the moment equalities. $\square$

[write] 6 October 2026; sources 15–20. *Credit.* This normal form is source 14’s Theorem II.44.4 (same freedom in the representatives c_x , the same additive and lower-order terms, the same counts (II.54.3)); proved independently. The intake checked both statements against each other on groups that neither package uses (no mismatch).

The representatives c_x need not depend additively on x , and the coordinate-omitting functions may vary arbitrarily from fibre to fibre. These freedoms are invisible to the mixed-value arrangement test. Omitting them would give an incorrect rigidity statement. This normal form classifies joint equality, not all extremizers of the product-property lower bound.

II.55 A geometry–amplification dichotomy

Source 15, delivered Section 11.

The quantitative results also provide an alternative that distinguishes two reasons an arrangement estimate can be useful: either collisions are substantially below their cubic cap, improving amplification, or most of the cubic fibre mass lies on nearly Cartesian fibres.

Theorem II.55.1 (Dichotomy with explicit weights). *Let $B \neq \emptyset$, $0 < \eta < 1$, and $0 < \varepsilon \leq 1/100$. At least one of the following alternatives holds:*

(i) $C \leq (1 - \eta)\Gamma$, and therefore

$$R_8 \geq (1 - \eta)^{-6} \frac{R_2^7}{\Gamma^6};$$

(ii) *the fibres satisfying $C_x \geq (1 - \varepsilon)b_x^3$ carry at least $(1 - \eta/\varepsilon)\Gamma$ of the total cubic mass $\sum_x b_x^3$. Every such nonempty fibre is within $18\varepsilon b_x$ points of a Cartesian coset, with the intersection and size bounds in (II.52.1).*

The lower bound in (ii) is informative when $\eta < \varepsilon$.

Proof. If (i)’s collision inequality holds, apply Theorem II.47.3. Otherwise, $\sum_x (b_x^3 - C_x) < \eta\Gamma$. A bad fibre with $C_x < (1 - \varepsilon)b_x^3$ contributes more than εb_x^3 to this deficit. Their total cubic mass is consequently at most $\eta\Gamma/\varepsilon$. On every remaining nonempty fibre, $Q(B_x) \geq C_x \geq (1 - \varepsilon)b_x^3$, so Theorem II.52.2 applies. $\square$

[write] 6 October 2026; sources 15–20. *Credit.* A weaker restatement of source 14’s cubic-weight split in Section II.43 (with $\tau(\theta)$ in place of 18ε) together with its Corollary II.44.2.

Taking $\eta = \varepsilon^2$ gives a particularly simple statement: either there is a factor $(1 - \varepsilon^2)^{-6}$ improvement over the cubic-cap bound, or at least $1 - \varepsilon$ of the cubic fibre mass has the stated Cartesian approximation. This mass is not point mass $\sum_x b_x$ and not uniform mass on the set of fibres. A claim about either of those would require an additional fibre-size comparison.

The theorem gives geometric information about the domain. It does not supply a quantitative repair of the labels to the exact multiadditive form. The latter is a separate research problem in Section V.60.

II.56 Vertex weights and other extensions

Source 15, delivered Section 12.

The moment method is naturally a method for nonnegative measures, not only for subsets. There are two different levels of generality.

II.56.1 Arbitrary cube measures

Let each $m_h : T \times V \rightarrow [0, \infty)$ be any nonnegative real measure, and define R_d as its additive $2d$ -tuple count, summed over h . The moment identity, Theorem II.47.3, the spectral stability estimates, and the side-vector profile inequality continue to hold verbatim. No vertex representation is needed for these statements. An arbitrary collection of measures does not, however, have a meaningful bound $C \leq \sum_x b_x^3$ unless suitable additional data and hypotheses are given.

II.56.2 Weights arising from vertices

Let $f : X \times T \rightarrow [0, 1]$, with ϕ defined on $\text{supp } f$. Give a supported cube the weight

$$w(y, h, x) = \prod_{\omega \in \{0,1\}^k} f(y + \omega h, x),$$

and define $m_h(x, v)$ by summing this weight over cubes of mixed value v . The arrangement count is now the sum of the products of the $2d$ cube weights. Put

$$b_x = \sum_y f(y, x), \quad \beta_x = b_x/M, \quad \beta = \mathbb{E}_x \beta_x.$$

The same normalizations for a_d, χ, ρ_3 apply.

Theorem II.56.1 (Weighted collision and coupling bounds). *For vertex-product weights in $[0, 1]$,*

$$C \leq \sum_x b_x^3. \tag{II.56.1}$$

Thus all density-sensitive amplification bounds remain valid. If the map on $\text{supp } f$ has PP_γ , then the configuration lower bound, the collision lower bound, and the coupled bounds (II.49.6)–(II.49.7) also remain valid, with these weighted densities and counts.

Proof. First discard the label-equality condition. For each pair of cubes, retain only the formal vertex factors at $y, y + h$, and z in the product $w(y, h, x)w(z, h, x)$. There are three distinct formal positions because $k \geq 1$, even if their actual values coincide. All remaining factors are at most one. Thus

$$w(y, h, x)w(z, h, x) \leq f(y, x)f(y + h, x)f(z, x).$$

Summing over h, y, z and using the bijection $h \leftrightarrow y + h$ gives b_x^3 , proving (II.56.1).

For the configuration induction, give a configuration the product of f over its 4^j formal vertices and impose the same face-respect conditions. At a fixed lower-dimensional configuration ζ , replace $\mathbf{1}_{E_\zeta}$ in the proof of Lemma II.49.3 by the product of its vertex weights at that level, set to zero where a lower face fails. The product property allows every such nonnegative weight. All sums and Jensen steps are unchanged. The zero-dimensional base is $\mathcal{F}_0 = f$, so it is valid for fractional values

as well. The injection from configurations to cube pairs preserves the selected 2^{j+1} formal vertex factors; discarding the others can only increase the weight. Hence it gives the required inequality of weighted counts.

For the pair bound, replace the reference weight in Theorem II.49.5 by

$$\tilde{w}_y(x) = w(y, h, x) m_h(x, D_h \phi(y, x)),$$

with zero value on unsupported cubes. Its sum over y, x is exactly C_h . The weighted quadruple output is bounded by the weighted $R_2(h)$ because the four additional reference-cube factors are at most one. The two averaging steps therefore give the same inequality (II.49.5). Combining it with the weighted collision lower bound and the unchanged moment interpolation proves the assertion. $\square$

Remark II.56.2 (Exact saturation forces hard vertex weights). For $k \geq 1$, equality in (II.56.1) implies that every $f(y, x)$ is zero or one. Indeed, in the cap proof consider the term with $h = 0$ and $z = y$. Its left-hand weight is $f(y, x)^{2^{k+1}}$, its retained right-hand weight is $f(y, x)^3$, and the label condition is automatically satisfied. These are strictly unequal for $0 < f(y, x) < 1$. Since every termwise difference is nonnegative, equality of the sums rules out all fractional values. The exact set classifications then apply. This observation alone gives no useful dimension-free quantitative rounding theorem for near equality; a single diagonal term can be tiny relative to the full cubic sum.

II.56.3 Which extensions require new work?

The group generalization above allows composite moduli, unequal coordinate sizes, noncyclic groups, and finite abelian value groups. It does not assert an analogous theorem for nonabelian value groups, where an alternating additive relation has no canonical interpretation. Nor does it automatically handle a requirement that all vertices of an entire arrangement be distinct. Weights on side vectors are compatible with the Fourier argument, but their normalizations and geometric caps must be tracked separately.

II.57 Source 17's chapter: polynomial-loss restriction at Section 9 [write]

Added in the fourth write, 6 October 2026.

The next seven sections are source 17, a manuscript on Gowers's random restriction in Section 9 (Lemmas 9.2–9.3 and Corollary 9.4): from many respected additive sixteen-tuples, pass to a subset on which almost all are respected. Source 17 replaces the long product of degree-one cosine filters by one Fejér kernel and keeps the exact second moment. *Logically the chapter follows source 04's Sections II.8–II.9*, which prove the same moment lifting and a complementary restriction theorem over prime fields; source 04's text was in the repository at this source's pin and is not cited by it. It is printed here because no number printed earlier may move. Its introduction is Section F.22; its Lean relation, audits, questions, conclusion and appendices are in Part V (Sections V.69–V.74). Its notation is in the table at the head of this Part.

What is proved, in one paragraph. Over a prime field, with $L = \lceil (m+1)/(\alpha\eta) \rceil$ and $N \geq 2L^m(2L-1)^{m-1}\beta^{-(m-1)}$, a set whose map respects $\alpha|B|^{m-1}$ balanced m -tuples has a subset on which a $1-\eta$ proportion is respected and which keeps $\alpha^m\eta^{m-1}|B|^{m-1}/(2(m+2)^m)$ tuples, for

weights and vector-valued maps too (Theorem II.61.2); at $m = 16$ the retention is $\alpha^{16}\eta^{15}/(2 \cdot 18^{16})$, in place of $(\alpha\eta/4)^{2^{19}}$. The moment lifting $G_{2r} \geq \gamma^{r-1}|B|^{2r-1}$ without a density loss is source 04's (M2)–(M5) (Proposition II.59.1). The best retention a scalar filter certificate can give is determined (Proposition II.62.2, sharpening source 04's Proposition II.9.6), and a random-linear-map construction treats arbitrary systems in fixed characteristic with a rank gap (Theorem II.63.5; new in this report). Against source 04's restriction theorem: a better constant for $d \leq 97$, a worse threshold; both are printed (note after Theorem II.61.2). Prime moduli only: composite moduli need a new argument (Section II.64).

Relation to the formal statements (checked at HEAD 0b4890a1c). `lemma_9_3` and `corollary_9_4` (Sections08_09.lean, lines 77, 91) hold for every N , with exact companions (`Proofs09Restriction.lean`, lines 376, 416) and the retention $(\alpha\eta/4)^{2^{19}}$; the definitions match this chapter's T_{16} , G_{16} and $1 - \eta$ exactly. The theorem here proves their prime case with an explicit threshold and a far better constant (note in Section F.22); it is not formalized.

[write] **7 October 2026; sources 47–53.** Source 48 (Section II.186) makes the theorem of this chapter constructive for $B \subseteq \mathbb{F}_p$ and scalar maps, with the same retention, and so answers source 17's Question 9 in its ambient-size form.

II.58 Notation and elementary counting facts

Source 17, delivered Section 2.

Let V and H be finite-dimensional vector spaces over a prime field $\mathbb{F}_p$. Write $N = |V|$. A map $\phi : V \rightarrow H$ need not be linear. Let $m = 2r \geq 4$, and write

$$\varepsilon = (\underbrace{1, \dots, 1}_r, \underbrace{-1, \dots, -1}_r) \in \mathbb{F}_p^m.$$

For $x = (x_1, \dots, x_m) \in V^m$, the notation $\varepsilon \cdot x = 0$ means $\sum_j \varepsilon_j x_j = 0$ in V . We call such a tuple *additive*. It is *respected* if also $\sum_j \varepsilon_j \phi(x_j) = 0$ in H .

For a weight $w : V \rightarrow [0, 1]$, define

$$T_m(w) = \sum_{\varepsilon \cdot x = 0} \prod_{j=1}^m w(x_j), \quad (\text{II.58.1})$$

$$G_m(w, \phi) = \sum_{\substack{\varepsilon \cdot x = 0 \\ \varepsilon \cdot \phi(x) = 0}} \prod_{j=1}^m w(x_j). \quad (\text{II.58.2})$$

For $w = \mathbf{1}_B$, abbreviate these quantities to $T_m(B)$ and $G_m(B, \phi)$. The word “proportion” in the conclusion always refers to $G_m(B', \phi)/T_m(B')$, not to normalization by $|B'|^{m-1}$. The theorem guarantees $T_m(B') > 0$.

Lemma II.58.1 (The free-variable bound). *If $b = \sum_x w(x)$, then*

$$0 \leq G_m(w, \phi) \leq T_m(w) \leq b^{m-1}.$$

In particular, $T_m(B) \leq |B|^{m-1}$.

Proof. After choosing $x_1, \dots, x_{m-1}$, the additive relation determines x_m uniquely because its coefficient is -1 . Its weight is at most one. Summing the product of the first $m-1$ weights gives b^{m-1} . The first two inequalities follow from nonnegativity and restriction of the summation set. $\square$

The parameter α in the main theorem therefore measures the number of respected tuples against the maximal scale b^{m-1} . It is not merely a lower bound for their fraction among whatever additive tuples happen to be present. This distinction is important in sparse or highly structured examples.

We use normalized averages over random parameters, but unnormalized counts of tuples. On $\mathbb{T} = \mathbb{R}/\mathbb{Z}$ let $e(t) = \exp(2\pi it)$. On $\mathbb{F}_p$ write $e_p(a) = \exp(2\pi ia/p)$, which is well-defined independently of the integer representative of a .

II.59 Moment lifting without a density loss

Source 17, delivered Section 3.

The following elementary result is useful independently of the restriction theorem. We prove it in an arbitrary finite abelian group so that the normalization is visible.

Proposition II.59.1 (Density-sensitive graph moment inequality). *Let A be a nonempty subset of a finite abelian group K , let $b = |A|$, and let*

$$E_r(A) = \#\{(a_1, \dots, a_{2r}) \in A^{2r} : a_1 + \dots + a_r = a_{r+1} + \dots + a_{2r}\}.$$

For each integer $r \geq 2$,

$$E_r(A) \geq \frac{E_2(A)^{r-1}}{b^{r-2}}. \quad (\text{II.59.1})$$

For $r > 2$, equality holds exactly when A is a coset of a subgroup of K .

Proof. Let $Q = |K|$, and use the unnormalized Fourier transform $\widehat{\mathbf{1}_A}(\chi) = \sum_{a \in A} \overline{\chi(a)}$. Character orthogonality gives

$$E_r(A) = \frac{1}{Q} \sum_{\chi \in \widehat{K}} |\widehat{\mathbf{1}_A}(\chi)|^{2r}, \quad b = \frac{1}{Q} \sum_{\chi} |\widehat{\mathbf{1}_A}(\chi)|^2.$$

For $r > 2$, apply Hölder to

$$z^4 = (z^2)^{(r-2)/(r-1)} (z^{2r})^{1/(r-1)}, \quad z \geq 0.$$

Using normalized sums yields $E_2(A) \leq b^{(r-2)/(r-1)} E_r(A)^{1/(r-1)}$. Raising to the power $r-1$ proves (II.59.1); for $r=2$ it is an identity.

For equality with $r > 2$, equality in Hölder says that all nonzero $|\widehat{\mathbf{1}_A}(\chi)|$ are the same. The trivial character has value b , so that common magnitude is b . Fix $a_0 \in A$ and put $J = \langle A - A \rangle$. Equality in the triangle inequality shows that $|\widehat{\mathbf{1}_A}(\chi)| = b$ exactly when χ is constant on A , equivalently χ annihilates J . Thus the Fourier support is $J^\perp$. Parseval gives

$$b = \frac{|J^\perp| b^2}{Q} = \frac{b^2}{|J|},$$

so $b = |J|$. Since $A \subseteq a_0 + J$, equality of cardinalities makes $A = a_0 + J$. Conversely, the Fourier transform of a subgroup coset has magnitude either zero or b , so all the inequalities are equalities. $\square$

[write] **6 October 2026; sources 15–20.** *Credit.* This is source 04's (M2) with its coset equality case (Theorem II.8.1), and Corollary II.59.2 is source 04's (M3) (Corollary II.8.2); the remark on $N^2|B|$ is source 04's (M4)–(M5), and source 04's (M1) with complex weights contains the remark on fractional weights.

Source 04's (M6) (Proposition II.8.3) is stronger. The arrangement analogues are source 14's Theorem II.44.1 and source 15's Theorem II.47.3. Proved independently; the source does not claim the inequality as new.

Apply this proposition to the graph $A = \{(x, \phi(x)) : x \in B\} \subseteq V \oplus H$. Its size is $|B|$, and $E_r(A) = G_{2r}(B, \phi)$.

Corollary II.59.2 (From quadruples to arbitrary even order). *If $G_4(B, \phi) \geq \gamma|B|^3$, with $0 < \gamma \leq 1$, then*

$$G_{2r}(B, \phi) \geq \gamma^{r-1}|B|^{2r-1} \quad (r \geq 2). \quad (\text{II.59.2})$$

Proof. Substitute $E_2(A) \geq \gamma|B|^3$ and $b = |B|$ into (II.59.1). The exponent of $|B|$ is $3(r-1) - (r-2) = 2r - 1$. $\square$

For the original absolute normalization $G_4(B, \phi) \geq \delta N^3$ and $|B| = \beta N$, the order-eight consequence is

$$G_{16}(B, \phi) \geq \delta^7 \beta^{-6} N^{15}. \quad (\text{II.59.3})$$

Putting $\delta = \gamma\beta^3$ gives precisely $\gamma^7 \beta^{15} N^{15}$, rather than $\gamma^7 \beta^{21} N^{15}$. The equality behind the improvement is the exact second moment $N^2|B|$, before it is bounded by N^3 . The inspected Lean proof contains that exact equality inside the proof of `section9_second_moment_1e`; it is therefore a natural point at which to expose a stronger reusable lemma [86].

Remark II.59.3 (Fractional graph weights). For a nonnegative graph weight f with values at most one, the same interpolation has denominator $(\sum f^2)^{r-2}$. Since $\sum f^2 \leq \sum f$, a hypothesis $E_2(f) \geq \gamma(\sum f)^3$ again implies $E_r(f) \geq \gamma^{r-1}(\sum f)^{2r-1}$. No equality classification for fractional weights is asserted here.

II.60 A single Fejér filter

Source 17, delivered Section 4.

II.60.1 The kernel and its signal

For an integer $L \geq 2$, define

$$K_L(t) = \frac{1}{L^2} \left| \sum_{a=0}^{L-1} e(at) \right|^2 = \frac{1}{L} \sum_{n=-(L-1)}^{L-1} \left(1 - \frac{|n|}{L}\right) e(nt). \quad (\text{II.60.1})$$

This normalization has $0 \leq K_L(t) \leq 1$, so it is an admissible inclusion probability. Its mean is $1/L$, and its Fourier coefficients are nonnegative.

Lemma II.60.1 (Exact relation signal). *For $m \geq 2$ let*

$$S_{m,L} = \sum_{n=-(L-1)}^{L-1} \left(1 - \frac{|n|}{L}\right)^m = 1 + 2 \sum_{j=1}^{L-1} \left(\frac{j}{L}\right)^m. \quad (\text{II.60.2})$$

Then

$$\frac{2L}{m+1} \leq S_{m,L} \leq L. \quad (\text{II.60.3})$$

For fixed m , $S_{m,L}/L \rightarrow 2/(m+1)$ as $L \rightarrow \infty$.

Proof. For the lower bound apply the trapezoidal inequality for the convex function $f(t) = (1-t/L)^m$ on $[0, L]$:

$$\frac{1}{2}f(0) + \sum_{j=1}^{L-1} f(j) + \frac{1}{2}f(L) \geq \int_0^L f(t) dt = \frac{L}{m+1}.$$

Twice the left side is $S_{m,L}$. For the upper bound use $u^m \leq u$ for $0 \leq u \leq 1$ and sum the triangular coefficients, whose sum is L . The limit is the corresponding Riemann-sum limit. $\square$

The signal in a one-layer degree-one filter is only $1 + 2^{1-m}$. In contrast, the signal (II.60.2) grows linearly with L . Choosing L proportional to $(\alpha\eta)^{-1}$ gives the whole needed amplification in one layer. Its cost is the controlled increase in the set of possible short Fourier relations.

II.60.2 Bounded affine degeneracy

The independent random translation used below is useful: it forces Fourier coefficient vectors to have coordinate sum zero. This saves one power in the elementary count of exceptional coefficient vectors. All coefficient sums in this definition are interpreted modulo p .

Definition II.60.2 (Exceptional tuples). Suppose $p > 2L - 2$. An additive tuple $x \in V^m$ is L -degenerate if there is an integer vector $n = (n_1, \dots, n_m) \in [-L+1, L-1]^m$ such that

$$\sum_j n_j = 0 \text{ in } \mathbb{F}_p, \quad \sum_j n_j x_j = 0 \text{ in } V,$$

and whose reduction modulo p does not belong to $\text{span}_{\mathbb{F}_p}(\varepsilon)$. A tuple that is not L -degenerate is L -generic.

Lemma II.60.3 (Exceptional-set bound). *The number of L -degenerate additive tuples in V^m is at most*

$$D_0 = (2L - 1)^{m-1} N^{m-2}. \quad (\text{II.60.4})$$

Every L -generic additive tuple has pairwise distinct entries.

Proof. There are at most $(2L - 1)^{m-1}$ vectors n under consideration with $\sum n_j = 0$ modulo p : after choosing the first $m - 1$ entries, the last entry is determined modulo p , and the interval of $2L - 1$ allowed integers contains at most one representative because $p > 2L - 2$. For an n outside $\text{span}(\varepsilon)$, the two equations with rows n and ε are independent over $\mathbb{F}_p$. They have exactly N^{m-2} solutions in V^m , as is seen by row reduction in each coordinate of V . A union bound proves (II.60.4).

If $x_i = x_j$ for $i \neq j$, the vector $e_i - e_j$ gives an additional allowed affine relation. It is not a multiple of ε , since $m \geq 4$ and every coordinate of ε is nonzero. Thus any repeated-entry tuple is degenerate. $\square$

Lemma II.60.4 (Short vectors on the principal relation line). *If $p > 2L - 2$ and $n \in [-L+1, L-1]^m$ reduces modulo p to a multiple of ε , then $n = t\varepsilon$ as an integer vector for a unique integer $t \in [-L+1, L-1]$.*

Proof. Take $t = n_1$, since $\varepsilon_1 = 1$. For each j , $n_j - \varepsilon_j t$ is divisible by p and has absolute value at most $2L - 2 < p$. It must therefore be zero. $\square$

II.60.3 The exact survival probabilities

Choose independently and uniformly

$$u \in V^*, \quad v \in H^*, \quad c \in \mathbb{F}_p.$$

Conditioned on these choices, include each $x \in V$ in a random subset B' independently with probability

$$w(x)K_L \left(\frac{u(x) + v(\phi(x)) + c}{p} \right). \quad (\text{II.60.5})$$

The notation denotes evaluation through e_p ; it does not require a preferred lift of the numerator. Points with $w(x) = 0$ are never chosen.

Lemma II.60.5 (Good and bad tuples separate exactly). *Set $a = L^{-m}$. For an L -generic additive tuple x , its probability of being contained in the random subset is*

$$\begin{cases} aS_{m,L} \prod_j w(x_j), & \varepsilon \cdot \phi(x) = 0, \\ a \prod_j w(x_j), & \varepsilon \cdot \phi(x) \neq 0. \end{cases} \quad (\text{II.60.6})$$

Proof. The entries are distinct by Lemma II.60.3; hence conditional independence really does give the product of the m inclusion probabilities. Expand all the factors using (II.60.1). Averaging over u, v, c leaves exactly the coefficient vectors n with

$$\sum_j n_j x_j = 0, \quad \sum_j n_j \phi(x_j) = 0, \quad \sum_j n_j = 0.$$

Indeed, a uniform linear functional on a finite vector space has average character zero on any nonzero vector and average one at zero. Genericity and Lemma II.60.4 force $n = t\varepsilon$ with $|t| \leq L - 1$. If the tuple is respected, every such t survives, giving $aS_{m,L}$. Otherwise its image discrepancy is a nonzero vector in H . Multiplication by a nonzero $t \in \mathbb{F}_p$ cannot annihilate that vector; thus only $t = 0$ survives, giving a . The fixed product of the weights multiplies both expressions. $\square$

II.61 The restriction theorem with all exceptional tuples included

Source 17, delivered Section 5.

It is not enough to obtain a large expected number of good tuples and a small expected number of bad ones separately. The same subset must satisfy both requirements. A single signed objective gives that simultaneity directly.

Proposition II.61.1 (An exact-data restriction certificate). *In the setting above, let $T = T_m(w)$, $G = G_m(w, \phi)$, and let D be any upper bound for the number of L -degenerate additive tuples in $(\text{supp } w)^m$. Let $0 < \eta \leq 1$. Define*

$$\Psi_L = L^{-m} [\eta S_{m,L} G - (1 - \eta)(T - G)] - D. \quad (\text{II.61.1})$$

If $\Psi_L > 0$, some $B' \subseteq \text{supp } w$ satisfies

$$G_m(B', \phi) \geq (1 - \eta)T_m(B'), \quad T_m(B') \geq \frac{\Psi_L}{\eta} > 0. \quad (\text{II.61.2})$$

Proof. For a realization of the random subset define X to be the number of respected additive tuples in it and Y the number of unrespected ones. Consider

$$F = \eta X - (1 - \eta)Y = X - (1 - \eta)(X + Y).$$

Let D_G and D_B be the numbers of degenerate respected and degenerate unrespected additive tuples in the initial support. Thus $D_G + D_B \leq D$. The total weighted mass of degenerate respected tuples is at most D_G . Lemma II.60.5 gives

$$\mathbb{E}F \geq \eta a S_{m,L}(G - D_G) - (1 - \eta)a(T - G) - (1 - \eta)D_B.$$

The last term bounds the entire possible contribution of degenerate bad tuples, including those with repeated entries. Since $a S_{m,L} \leq 1$ by Lemma II.60.1, the two exceptional penalties are at most $D_G + D_B \leq D$. Therefore $\mathbb{E}F \geq \Psi_L$. There is a realization with $F \geq \Psi_L$. For that realization $F \geq 0$ implies $X \geq (1 - \eta)(X + Y)$, and $F \leq \eta(X + Y)$ implies $X + Y \geq \Psi_L/\eta$. These are exactly (II.61.2). $\square$

The proposition is stronger than a worst-case threshold statement: for a particular input one may count or estimate its actual exceptional tuples and optimize over L . No union-bound replacement is compulsory.

Theorem II.61.2 (Weighted, vector-valued polynomial restriction). *Let $m = 2r \geq 4$. Let V, H be finite-dimensional $\mathbb{F}_p$ -vector spaces, $N = |V|$, $w : V \rightarrow [0, 1]$, and $\phi : V \rightarrow H$. Write $b = \sum_x w(x) > 0$ and $\beta = b/N$. Suppose $0 < \alpha, \eta \leq 1$ and*

$$G_m(w, \phi) \geq \alpha b^{m-1}. \quad (\text{II.61.3})$$

Let $L = \lceil (m+1)/(\alpha\eta) \rceil$. If

$$p > 2L - 2, \quad N \geq 2L^m(2L - 1)^{m-1}\beta^{-(m-1)}, \quad (\text{II.61.4})$$

there is a subset $B' \subseteq \text{supp } w$ such that

$$G_m(B', \phi) \geq (1 - \eta)T_m(B'), \quad (\text{II.61.5})$$

$$T_m(B') \geq \frac{b^{m-1}}{2\eta L^m} \geq \frac{\alpha^m \eta^{m-1}}{2(m+2)^m} b^{m-1}. \quad (\text{II.61.6})$$

Proof. Put $M = b^{m-1}$ and $a = L^{-m}$. By Lemma II.58.1, $T \leq M$. By Lemma II.60.1 and the choice of L ,

$$\alpha\eta S_{m,L} \geq \frac{2\alpha\eta L}{m+1} \geq 2.$$

The coefficient of G in (II.61.1) is positive, so using $G \geq \alpha M$ and $T \leq M$ gives

$$\begin{aligned} a[\eta S_{m,L}G - (1 - \eta)(T - G)] &\geq a[\alpha\eta S_{m,L} - (1 - \eta)(1 - \alpha)]M \\ &\geq aM. \end{aligned}$$

The ambient-size hypothesis is precisely the sufficient inequality

$$D_0 = (2L - 1)^{m-1}N^{m-2} \leq \frac{1}{2}a\beta^{m-1}N^{m-1} = \frac{1}{2}aM.$$

Thus $\Psi_L \geq aM/2 > 0$. Proposition II.61.1 gives (II.61.5) and the first part of (II.61.6). Finally, $\alpha\eta \leq 1$ implies

$$L \leq \frac{m+1}{\alpha\eta} + 1 \leq \frac{m+2}{\alpha\eta},$$

which yields the second lower bound. $\square$

Proof of Theorem F.22.1. Take $V = H = \mathbb{F}_N$ and $w = \mathbf{1}_B$, extending ϕ arbitrarily outside B . The size threshold implies the characteristic condition. The lower bound for $|B'|$ follows from $T_m(B') \leq |B'|^{m-1}$. $\square$

[write] **6 October 2026; sources 15–20.** *Against source 04's restriction theorem* (Theorem II.9.1; checked at the write). The normalizations agree ($m = 2d$, ordered tuples with repetitions, $\alpha|B|^{m-1}$), and so do the powers $\alpha^m \eta^{m-1}$. *Constants:* the retention $1/(2(m+2)^m) = 1/(2(4(d+1)^2)^d)$ here against source 04's $\kappa_d = 1/(2(400d)^d)$; this one is larger exactly when $(d+1)^2 < 100d$, that is for $d \leq 97$ (at $d = 8$ by the factor $3200^8/18^{16} \approx 9.05 \cdot 10^7$). *Thresholds:* (II.61.4) has degree $2m-1$ in $(\alpha\eta)^{-1}$ and $\beta^{-(m-1)}$, source 04's (R4) degree m and β^{-1} , far better. At the Corollary 9.4 point $\gamma = \eta = \beta = 1/2$ (relative to $\beta^{15}N^{15}$): Corollary II.61.3 retains $2^{-194.72}$ above $N \geq 2^{405.71}$; source 04's Corollary II.9.5 retains 2^{-222} above 2^{231} ; `corollary_9_4` retains $2^{-8,388,608}$. Source 04's conclusion ratio $1/(1+\eta)$ is slightly stronger than $1-\eta$; this theorem's weights and vector-valued maps over $\mathbb{F}_p$ -spaces are broader than source 04's $B \subseteq \mathbb{F}_p$. Both are printed; neither dominates. A composition of the two (this filter with source 04's per-tuple lower bound and weighted rank-two count) is a marked question of the intake, not proved here (item II.6).

[write] **7 October 2026; sources 47–53.** *Made constructive by source 48* (Section II.186) in the unweighted scalar case $B \subseteq \mathbb{F}_p$: its Theorem II.194.1 constructs B' with the same retention $b^{m-1}/(2\eta L^m)$ by rational arithmetic in $O(C_s n p^5)$ operations, through exact collision-aware energy inventories; this answers this source's Question 9 in its ambient-size form (item II.23).

[write] **7 October 2026; sources 60–69.** *Threshold lowered by source 60* (Section II.310): with the retention $\alpha^m \eta^{m-1}/(2(m+2)^m)$ kept, the size condition $b \geq 2(1 + \binom{m}{2})L^{m-1}$, $L = \lceil (m+1)/(\alpha\eta) \rceil$, suffices on prime moduli (its Corollary II.330.5); this proves the intake's marked composition with one power of L less (item II.6).

Corollary II.61.3 (A density-sensitive replacement for Corollary 9.4). *Let N be prime, $|B| = \beta N > 0$, and suppose $G_4(B, \phi) \geq \gamma \beta^3 N^3$ with $0 < \gamma \leq 1$. Let $0 < \eta \leq 1$ and $L = \lceil 17/(\gamma^7 \eta) \rceil$. If*

$$N \geq 2L^{16}(2L-1)^{15}\beta^{-15},$$

there exists $B' \subseteq B$ such that

$$G_{16}(B', \phi) \geq (1-\eta)T_{16}(B'), \quad T_{16}(B') \geq \frac{\gamma^{112}\eta^{15}}{2 \cdot 18^{16}}|B|^{15}.$$

Proof. Corollary II.59.2 with $r = 8$ gives $G_{16}(B, \phi) \geq \gamma^7|B|^{15}$. Apply Theorem F.22.1 with $m = 16$ and $\alpha = \gamma^7$. $\square$

II.61.1 A quantitative comparison

Table II.61.1 uses $\beta = 1/2$, $m = 16$, and the exact integer $L = \lceil 17/(\alpha\eta) \rceil$. The old and new retention factors multiply $|B|^{15}$. The last column is the base-ten logarithm of the sufficient threshold in (F.22.2), not a claim of its optimality. The data are generated by the accompanying script.

α	η	L	$\log_{10}$ old	$\log_{10}$ new	$\log_{10} N_0$
1/2	1/10	340	−997767.24	−39.80	87.80
1/10	1/10	1700	−1364228.83	−50.99	109.47
1/100	1/100	170000	−2412804.83	−81.99	171.48

Table II.61.1: Local retention factors and the explicit sufficient ambient threshold.

The new estimates remain far too large for a practical application at these sample parameters. Their point is the change in parameter dependence: a polynomial degree tied to the number of variables replaces a very large fixed exponent. Both the retained mass and the exceptional set must be considered when assessing the quantitative gain.

II.62 A limitation of the scalar filter-certificate method

Source 17, delivered Section 6.

The polynomial exponents have an intrinsic explanation within the particular expectation argument used here. The following does not give an extremal theorem about all possible subsets.

Lemma II.62.1 (Universal good-tuple moment bound). *Let $h : \mathbb{T} \rightarrow [0, 1]$ be a trigonometric polynomial, with Fourier coefficients c_n and mean $q = c_0 > 0$. For even $m \geq 4$ put $P_g = \sum_n |c_n|^m$. Then*

$$P_g \leq q^{m-1}. \quad (\text{II.62.1})$$

In the generic balanced-relation model, the bad-tuple survival is $P_b = q^m$, so its distinguishing ratio satisfies $P_g/P_b \leq 1/q$.

Proof. Since $h \geq 0$, $|c_n| \leq \int h = q$. Parseval and $h^2 \leq h$ give $\sum_n |c_n|^2 \leq q$. Therefore

$$\sum_n |c_n|^m \leq q^{m-2} \sum_n |c_n|^2 \leq q^{m-1}.$$

For a generic balanced relation the only Fourier vectors are $n\varepsilon$; their coefficients multiply to $c_n^{m/2} c_{-n}^{m/2} = |c_n|^m$. A generic unrespected tuple retains only the zero vector, giving q^m . $\square$

The same inequality holds for a function on a finite abelian group with normalized Fourier transform. For the indicator of a subgroup, it is an equality: all nonzero Fourier coefficients have magnitude q . This explains why the random-linear-map model of the next section can attain this moment bound exactly.

Proposition II.62.2 (Best possible powers for this scalar certificate). *Consider the ideal model with no exceptional tuples, total initial count T , and respected fraction $\rho = G/T$. Assume $0 < \eta < 1$ and $0 < \rho < 1 - \eta$. A certificate using only the exact generic survival probabilities of one filter and the signed objective in Proposition II.61.1 cannot certify a retained count larger than*

$$\frac{(m-1)^{m-1}}{m^m} \frac{\rho^m \eta^{m-1}}{(1-\eta)^{m-1} (1-\rho)^{m-1}} T. \quad (\text{II.62.2})$$

The same upper bound applies to any finite product of independent filters within this scalar-certificate model.

Proof. Write $c = (1 - \eta)/\eta$. The count certified by the signed expectation, divided by T , is

$$\rho P_g - c(1 - \rho) P_b.$$

Lemma II.62.1 bounds this by

$$f(q) = \rho q^{m-1} - c(1 - \rho) q^m.$$

Its maximum on $[0, 1]$ occurs at

$$q_* = \frac{(m-1)\rho}{mc(1-\rho)}.$$

The hypothesis $\rho < 1 - \eta$ ensures $q_* < 1$. Substitution gives (II.62.2). For independent filters with means q_i , the bad probability is $(\prod_i q_i)^m$ and the good probability is at most $(\prod_i q_i)^{m-1}$. The same calculation applies with $q = \prod_i q_i$. $\square$

[write] **6 October 2026; sources 15–20.** *Credit.* Lemma II.62.1 is the inequality of source 04’s Proposition II.9.6 ($Q_K \leq 1/b_0$ is $P_g \leq q^{m-1}$); this proposition sharpens it to an unconditional optimum of the scalar certificate in terms of the respected fraction ρ , where source 04 has only the conditional form “if $Q_K \geq 4/(\alpha\eta)$ ”.

Remark II.62.3 (The precise scope of the limitation). For small ρ and η , the displayed upper bound has powers $\rho^m \eta^{m-1}$. The Fejér argument achieves these powers, up to constants depending on m and the exceptional-set requirement. This does *not* prove that no subset can be larger, that no more informative probability argument can do better, or that the main theorem is globally sharp. In particular, using more information about the input than G, T can evade this limitation.

There is a second, simpler constraint: a filter of degree D has at most $2D + 1$ Fourier coefficients, and $|c_n| \leq q$. Hence its distinguishing ratio is at most $2D + 1$. A one-stage ratio of order $(\alpha\eta)^{-1}$ needs degree of that order. The present choice $D = L - 1$ has the correct order for fixed m .

II.63 Fixed characteristic: arbitrary linear systems and rank profiles

Source 17, delivered Section 7.

The Fejér theorem requires large characteristic. A genuinely different construction gives an analogous restriction theorem over a fixed field when the domain dimension is large enough. It also exposes a useful rank parameter that is invisible for a single relation.

II.63.1 The relation space

Let $\mathcal{R} \subseteq \mathbb{F}_p^m$ have dimension d , where $1 \leq d \leq m - 2$, and assume

$$\mathcal{R} \subseteq \mathbf{1}^\perp = \{a \in \mathbb{F}_p^m : \sum_j a_j = 0\}. \quad (\text{II.63.1})$$

Thus the system is invariant under translating all entries by the same vector. Put $s = m - d \geq 2$. A tuple $x \in V^m$ is a solution if

$$\sum_j a_j x_j = 0 \quad \text{for every } a \in \mathcal{R}.$$

It is respected by ϕ if the same equations hold with $\phi(x_j)$. Let $T_{\mathcal{R}}(B)$ and $G_{\mathcal{R}}(B, \phi)$ count these tuples. Row reduction shows

$$T_{\mathcal{R}}(B) \leq |B|^s. \quad (\text{II.63.2})$$

Indeed, d pivot coordinates are determined by the other s coordinates.

Definition II.63.1 (Affine genericity and error rank). A solution tuple is affine-generic if

$$\left\{ a \in \mathbb{F}_p^m : \sum_j a_j = 0, \sum_j a_j x_j = 0 \right\} = \mathcal{R}. \quad (\text{II.63.3})$$

For a generic solution, its *error rank* is the rank of the map

$$\mathcal{R} \longrightarrow H, \quad a \longmapsto \sum_j a_j \phi(x_j). \quad (\text{II.63.4})$$

Write this rank as $j(x)$. It is zero exactly for respected solutions.

Lemma II.63.2 (Degeneracy of a system). *The number of nongeneric solution tuples in V^m is at most*

$$D_{\mathcal{R}} = C_{s,p} N^{s-1}, \quad C_{s,p} = \frac{p^{s-1} - 1}{p - 1}. \quad (\text{II.63.5})$$

Proof. An extra affine relation is represented by a nonzero element of $\mathbf{1}^\perp/\mathcal{R}$, a vector space of dimension $s - 1$. Relations differing by a nonzero scalar or by an element of $\mathcal{R}$ impose the same additional equation on solution tuples. There are $(p^{s-1} - 1)/(p - 1)$ projective classes. For each representative, adjoining its row increases the rank by one, leaving exactly N^{s-1} solutions in V^m . Take a union bound. $\square$

II.63.2 Random linear maps see the exact error rank

Choose a uniform linear map

$$A : V \oplus H \oplus \mathbb{F}_p \longrightarrow \mathbb{F}_p^k,$$

and let

$$B_A = \{x \in B : A(x, \phi(x), 1) = 0\}. \quad (\text{II.63.6})$$

Set $q = p^{-k}$.

Lemma II.63.3 (Exact rank-profile survival). *For an affine-generic solution x of error rank j ,*

$$\Pr[x_1, \dots, x_m \in B_A] = q^{s+j}. \quad (\text{II.63.7})$$

This formula does not require distinct entries or independent Bernoulli choices at the points.

Proof. Consider the augmented column vectors $z_i = (x_i, \phi(x_i), 1)$. The relation space among the projected columns $(x_i, 1)$ is $\mathcal{R}$ by genericity. The relation space among the full z_i is the kernel of (II.63.4) inside $\mathcal{R}$, of dimension $d - j$. Thus the span of the z_i has dimension $m - (d - j) = s + j$. A uniform linear map to $\mathbb{F}_p^k$ kills a fixed subspace of dimension $s + j$ with probability $p^{-k(s+j)}$. This is exactly the event that every selected point belongs to B_A . $\square$

Proposition II.63.4 (A rank-profile certificate). *Let n_j be the number of affine-generic solution tuples in B^m with error rank j , and let D bound the nongeneric solution tuples. If*

$$\Omega(q) = \eta q^s n_0 - (1 - \eta) \sum_{j \geq 1} q^{s+j} n_j - D > 0, \quad (\text{II.63.8})$$

some set B_A has respected proportion at least $1 - \eta$ and

$$T_{\mathcal{R}}(B_A) \geq \Omega(q)/\eta.$$

Proof. Apply Lemma II.63.3 to the signed objective $G_{\mathcal{R}}(B_A, \phi) - (1 - \eta)T_{\mathcal{R}}(B_A)$. The possible negative contribution of nongeneric tuples is at least $-D$. The expectation is therefore at least $\Omega(q)$. The same deterministic inequalities used in Proposition II.61.1 finish the proof. $\square$

The profile matters: a tuple breaking several independent relations is suppressed by several extra powers of q . Collapsing all bad tuples into a single class throws away this information.

Theorem II.63.5 (Restriction with a rank gap). *Let $|B| = \beta N > 0$ and suppose*

$$G_{\mathcal{R}}(B, \phi) \geq \rho|B|^s, \quad 0 < \rho, \eta \leq 1.$$

Assume that every affine-generic unrespected solution has error rank at least h , where $1 \leq h \leq d$. Choose the least positive integer k for which, with $q = p^{-k}$,

$$q^h \leq \rho\eta/2. \quad (\text{II.63.9})$$

If

$$N \geq \frac{8C_{s,p}}{\rho\eta q^s \beta^s}, \quad (\text{II.63.10})$$

there is $B' \subseteq B$ with respected proportion at least $1 - \eta$ and

$$T_{\mathcal{R}}(B') \geq \frac{\rho q^s}{4} |B|^s \geq \frac{\rho^{1+s/h} \eta^{s/h}}{4 \cdot 2^{s/h} p^s} |B|^s. \quad (\text{II.63.11})$$

In particular the theorem always applies with $h = 1$ when its size hypothesis holds.

Proof. Write $M = |B|^s$ and $D_0 = C_{s,p}N^{s-1}$. There are at least $\rho M - D_0$ generic respected tuples, and at most M generic unrespected tuples in total by (II.63.2). The latter survive with probability at most q^{s+h} . For the signed objective F we therefore have

$$\begin{aligned} \mathbb{E}F &\geq \eta q^s (\rho M - D_0) - (1 - \eta) q^{s+h} M - D_0 \\ &\geq q^s (\eta \rho - q^h) M - 2D_0 \\ &\geq \frac{1}{2} \rho \eta q^s M - 2D_0. \end{aligned}$$

Condition (II.63.10) says $D_0 \leq \rho \eta q^s M / 8$, so $\mathbb{E}F \geq \rho \eta q^s M / 4$. Choose a realization at least its expectation and divide the resulting count bound by η . Minimality of k gives $q > p^{-1}(\rho\eta/2)^{1/h}$, including when $k = 1$. This proves the second inequality in (II.63.11) after harmless weakening from strict to non-strict inequality. Every unrespected generic solution has positive error rank, so $h = 1$ requires no additional hypothesis. $\square$

[write] **6 October 2026; sources 28–33.** Source 29's Section II.120 keeps the affine-rank stratification of every tuple, degenerate ones included, which removes the ambient-dimension requirement and makes the threshold linear in $(\alpha\eta)^{-1}$ (Theorems II.120.4, II.120.6); that answers Question 2 of this source for fixed fields (item II.6).

Corollary II.63.6 (One balanced relation in fixed characteristic). *For $m = 2r \geq 4$, take $\mathcal{R} = \text{span}(\varepsilon)$, so $d = 1$ and $s = m - 1$. Under the size hypothesis of Theorem II.63.5, $G_m(B, \phi) \geq \rho|B|^{m-1}$ yields a subset with respected proportion at least $1 - \eta$ and*

$$T_m(B') \geq \frac{\rho^m \eta^{m-1}}{4(2p)^{m-1}} |B|^{m-1}.$$

II.63.3 Why this does not replace the prime-cyclic proof

A generic system tuple requires the augmented vectors $(x_i, 1)$ to have rank s . Necessarily $\dim V + 1 \geq s$. For the sixteen-variable single relation, $s = 15$; one needs $\dim V \geq 14$ before a generic tuple can even exist. When $V = \mathbb{F}_N$ is one-dimensional, all such tuples are degenerate in the full-field sense. The Fejér construction succeeds there because it only excludes *bounded-coefficient* extra relations, not every relation over the field. This distinction is essential to the applicability of the two theorems.

II.64 Boundary cases, normalization traps, and constructive content

Source 17, delivered Section 8.

II.64.1 Repeated entries really are different events

For a repeated tuple, multiplying the point-inclusion probability once per occurrence is incorrect. For example, with $L = 2$ and a random translation c over $\mathbb{F}_7$, the event that the tuple $(0, 0, 0, 0)$ is contained in the selected subset has probability

$$\mathbb{E}_c K_2(c/7) = \frac{1}{2}.$$

The formal product of four identical probabilities instead gives

$$\mathbb{E}_c K_2(c/7)^4 = \frac{35}{128}.$$

Our product formula is only asserted for generic tuples, which have distinct entries. [6 October 2026; sources 15–20] *Precision:* the product with multiplicities is wrong as an identity but valid as a lower bound, since $0 \leq K_L \leq 1$ (here $35/128 \leq 1/2$); a lower bound is all a respected-tuple count needs, and source 04’s survival lemma (Lemma II.9.4) counts all respected tuples, repeated ones included. The final theorem then pays for every exceptional tuple, good or bad, so its conclusion is about the original tuple count.

II.64.2 Large characteristic cannot be omitted from the filter formula

For $p = 3$ and $L = 4$ one has

$$K_4(0) = 1, \quad K_4(1/3) = K_4(2/3) = \frac{1}{16}.$$

The discrete mean is $3/8$, not $1/L = 1/4$. Fourier coefficients have aliased modulo p . This example refutes the unrestricted use of the intermediate formula, not the possibility of a different small-characteristic restriction theorem. Section II.63 supplies one such theorem with different hypotheses.

II.64.3 Composite cyclic groups require a separate argument

In $\mathbb{Z}/N\mathbb{Z}$ with composite N , a nonzero scalar may annihilate a nonzero image discrepancy. Also two formally distinct coefficient rows need not define a map with fibers of size N^{m-2} . Both issues enter

the present proof, respectively in Lemma II.60.5 and Lemma II.60.3. The prime-cyclic improvement should therefore not be substituted for an all-moduli proposition without a new torsion analysis. [6 October 2026; sources 15–20] *Two precisions* (HEAD 0b4890a1c). The formal all- N proof avoids the issue differently: its height-zero restriction uses degree-one filters with digit coefficients, whose new relations have a coefficient in $\{\pm 1, \pm 2\}$ with fibres of size at most two in every cyclic group (Proofs15Restriction.lean, lines 2083–2086). And the obstruction needs a small prime factor: the intake observed that Lemmas II.60.3–II.60.5 carry over verbatim when every prime factor of N exceeds $2L - 2$ (the relevant minors are integers of absolute value at most $2L - 2$, and every $0 < |t| < L$ is a unit); that observation is not written in the source and is recorded as a marked question (item II.8).

II.64.4 Trivial successful inputs

If $G_m(B, \phi) \geq (1 - \eta)T_m(B)$ already, take $B' = B$; no threshold or filter is necessary. Likewise $\eta = 1$ imposes no quality restriction. The uniform theorem remains true at $\eta = 1$, but is deliberately not optimized for that trivial endpoint. If B is a subgroup coset and ϕ is affine on it, every tuple is respected. The small retention lower bound is not a description of the actual best subset on such an input.

II.64.5 Existence, derandomization, and running time

All random objects here take values in finite probability spaces apart from the notation for the Bernoulli probabilities. The proof gives a finite weighted average over subsets and parameter choices, so a subset with the required signed objective exists. For the Fejér construction, one can enumerate u, v, c and then fix the Bernoulli choices one at a time by conditional expectation. For the linear-map construction, enumeration of the finite matrix space is enough.

These are finite constructive procedures, not efficient algorithms as stated. A direct evaluation of the objective may enumerate $|B|^{m-1}$ tuples, and the full parameter space can itself be very large. Positive expected surplus does not by itself imply that a satisfactory subset is found with a substantial probability. No polynomial-time claim in $\log N$, or practical sampling guarantee, is made.

II.65 Sources 18 and 19: degenerate arrangements (Section 15) and common sampling (Section 16) [write]

Added in the fourth write, 6 October 2026.

The next six sections treat two later stages of Gowers’s inverse argument that no earlier source of this report touches. *Source 18’s Section II.66* is on Section 15: Lemma 15.4 bounds the number of degenerate arrangements, and Lemma 15.5 needs that bound to absorb them when a Riesz product selects a denser respected set. It continues the arrangement chain from Lemma 14.8 that source 04’s Section II.10, source 14’s Section II.44 and source 15’s chapter (Section II.45) hold. *Source 19’s chapter* (Sections II.67–II.70 and its worked comparisons, Section II.71) is on Section 16: Lemma 16.10 samples cross-sections, interpolates the affine pieces through the sampled points and refines the resulting partitions. Both are printed here, at the end of Part II, because no number printed earlier may move. Source 18’s introduction is Section F.23, source 19’s Section F.24; their

other sections are in Parts I, III and V. Their notation is in the table at the head of this Part.

What is proved, in one paragraph. Source 18: for $d \geq 2$ and primes $p > 2d$ the proportion of degenerate d -arrangements is $(k + C_d)/p + O(p^{-2})$ with $C_d = (T_{2d} - 2\binom{2d}{d} + 1)/2$, sharp in its first-order constant (Theorem II.66.8; $C_8 = 2,585,444$), and Lemma 15.5 needs only $\lambda p \geq k + 2,585,445$ in place of $\lambda p \geq k3^{2^{k+4}}$ (Corollary II.66.9); the selection exponent is unchanged and nothing is propagated. Source 19: a common sample of $\min\{n, \lceil qd/\sigma \rceil - 1\}$ distinct columns loses at most σ of the mass (Theorem II.67.4), so order q/σ anchors suffice where the paper takes $q\sigma^{-2}$; the finite minimax value is exact (Theorem II.68.2) and $r\mathcal{M}/q \rightarrow C_d$, with $C_2 = \varphi^3 e^{-\varphi} = 0.83996\dots$ (Theorem II.69.4); refining the sampled slices once replaces a quadratic number of refinement steps by a linear one, with $\sum_j e_j(M_1, \dots, M_r)$ candidates (Theorem II.70.3); and a local replacement for the step in Lemma 16.10 (Corollary II.70.8), which needs an inherited-coordinate hypothesis. Lemma 16.10 itself is not proved: both displayed closing comparisons fail (Section V.80.2), as the source says.

Relation to the formal statements (checked at HEAD 0b4890a1c). `lemma_15_4` and `lemma_15_5` (Sections14_15.lean, lines 402, 414) have exact companions; the degeneracy predicate is the formal one; source 18's theorem implies the formal bound for $d \geq 2, p > 2d$ (notes there). `lemma_16_9` is an exact companion since 16bbe5ad, after source 19's pin; `lemma_16_10` (Section16.lean, line 851) is open. Formal results written after that pin confirm source 19's audit points (the distinct-anchor counterexample, the failure of the same-dimension comparison) and prove a local affine lift with weaker constants (notes after Theorem II.67.4 and Corollary II.70.8); they are not source 19's statements. Nothing in these sections is formalized.

II.66 A sharp first-order count of degenerate arrangements

Source 18, delivered Section 2.

The purpose of this section is to refine Lemma 15.4 of Gowers [1]. We retain the full definition of degeneracy given immediately before that lemma. This qualification matters: Gowers observes after his proof that several tests involving the last coordinate can be discarded. Those tests are exactly what produces the additional saving proved below. The result therefore applies to the original definition, and to the random selection in Lemma 15.5, which uses all the coordinate monomials. It is not a corresponding claim about the weaker definition mentioned in that remark.

For fixed cube dimension k and arrangement order $d \geq 2$, we determine the leading coefficient in the proportion of degenerate arrangements:

$$\frac{k + C_d}{p} + O_{k,d}(p^{-2}), \quad C_d = \frac{T_{2d} - 2\binom{2d}{d} + 1}{2}, \quad T_n = [z^0](z^{-1} + 1 + z)^n.$$

The coefficient $k + C_d$ is sharp. We also give a finite inequality, so that the asymptotic statement can be used without an unspecified error constant. In the case $d = 8$ relevant to Gowers's proof, $C_8 = 2,585,444$. Under the usual restrictions that the density and error parameters lie in $(0, 1]$, this replaces the factor $k3^{2^{k+4}}$ in the size hypothesis of Lemma 15.5 by $k + 2,585,445$. The random-selection exponent itself is unchanged.

These are refinements proved in this article. We make no claim that their precise formulation is absent from all previous literature. The numbered source statements refer to Gowers's paper and, where explicit hypotheses have been supplied, to the repository transcription [1, 88].

II.66.1 Arrangements and their certificates

Fix integers $k \geq 1$ and $d \geq 2$, and write

$$q = 2d, \quad s = (\underbrace{1, \dots, 1}_d, \underbrace{-1, \dots, -1}_d) \in \mathbb{F}_p^q, \quad H = \{r \in \mathbb{F}_p^q : s \cdot r = 0\}.$$

An arrangement is specified by

$$h \in \mathbb{F}_p^k, \quad r \in H, \quad y = (y^1, \dots, y^q) \in (\mathbb{F}_p^k)^q,$$

through the map

$$\rho(\epsilon, j) = (y_1^j + \epsilon_1 h_1, \dots, y_k^j + \epsilon_k h_k, r_j), \quad \epsilon \in \{0, 1\}^k, \quad 1 \leq j \leq q. \quad (\text{II.66.1})$$

Put

$$M = q2^k, \quad D = (q + 1)k + q - 1.$$

There are p^D arrangements. All probabilities in this section use independent uniform choices of h , $r \in H$, and y . Set

$$\pi(\epsilon) = (-1)^{\epsilon_1 + \dots + \epsilon_k}, \quad \eta_0(\epsilon, j) = s_j \pi(\epsilon).$$

Definition II.66.1 (Certificate and degeneracy). A *certificate* is a function $\eta : \{0, 1\}^k \times \{1, \dots, q\} \rightarrow \{-1, 0, 1\} \subset \mathbb{F}_p$. It is *successful* for ρ if

$$\sum_{\epsilon, j} \eta(\epsilon, j) \prod_{i \in A} \rho(\epsilon, j)_i = 0 \quad \text{for every } A \subseteq \{1, \dots, k + 1\}. \quad (\text{II.66.2})$$

The arrangement is *degenerate* if a successful certificate lies outside $\text{span}_{\mathbb{F}_p} \{\eta_0\}$. A certificate is *pure parity* if it has the form $\eta(\epsilon, j) = c_j \pi(\epsilon)$ for some $c \in \mathbb{F}_p^q$.

The certificate η_0 is successful for every arrangement. Indeed, summation against parity annihilates every proper horizontal coordinate monomial. The two remaining expressions are multiples of $\sum_j s_j = 0$ and $\sum_j s_j r_j = 0$. If one sidelength is zero, two vertices in a cube coincide. Assigning coefficients 1 and -1 to those vertices, and zero elsewhere, gives a successful certificate outside the span of η_0 . Consequently every such arrangement is degenerate, and their proportion is exactly

$$z_k(p) = 1 - (1 - p^{-1})^k. \quad (\text{II.66.3})$$

II.66.2 The rank supplied by the first nonzero moment

The next lemma is the main algebraic step. It uses the equations with the last coordinate to distinguish most cross-sections from the exceptional diagonal $r_1 = \dots = r_q$.

Lemma II.66.2 (Rank of a nonparity certificate). *Let p be an odd prime, and fix nonzero sidelengths $h_1, \dots, h_k$. For a nonparity certificate, define its Boolean moments by*

$$c_{j,S} = \sum_{\epsilon \in \{0,1\}^k} \eta(\epsilon, j) \prod_{i \in S} \epsilon_i.$$

Let $m < k$ be the smallest cardinality of a set S for which some $c_{j,S}$ is nonzero. Choose $B \subseteq [k]$ with $|B| = m$ and $b = (c_{1,B}, \dots, c_{q,B}) \neq 0$, and put $t = k - m$. Then the conditional probability that this certificate is successful is at most

$$\begin{cases} p^{-(t+1)}, & b \notin \text{span}\{s\}, \\ (1 - p^{-(q-2)})p^{-2t} + p^{-(q-2+t)}, & b \in \text{span}\{s\}. \end{cases} \quad (\text{II.66.4})$$

In particular, every nonparity certificate has conditional success probability at most

$$Q_d(p) = p^{-2} + (p - 1)p^{-2d}. \quad (\text{II.66.5})$$

Proof. First, m exists and is less than k . If ϵ_T denotes the Boolean vector with support T , the moment transform is

$$c_{j,S} = \sum_{T \supseteq S} \eta(\epsilon_T, j), \quad \eta(\epsilon_T, j) = \sum_{S \supseteq T} (-1)^{|S \setminus T|} c_{j,S}.$$

Thus vanishing of every proper moment in a slice is equivalent to that slice being a multiple of parity. If it happened in every slice, the certificate would be pure parity, contrary to the hypothesis.

For $A \subseteq [k]$, write $h_A = \prod_{i \in A} h_i$ and

$$P_{j,A}(y^j) = \sum_{\epsilon} \eta(\epsilon, j) \prod_{i \in A} (y_i^j + \epsilon_i h_i).$$

Expanding gives the exact identity

$$P_{j,A}(y^j) = \sum_{S \subseteq A} h_S c_{j,S} \prod_{i \in A \setminus S} y_i^j. \quad (\text{II.66.6})$$

All moments of degree less than m vanish. Hence $P_{j,B} = h_B b_j$. The tests indexed by B and by $B \cup \{k+1\}$ therefore imply

$$\sum_j b_j = 0, \quad b \cdot r = 0. \quad (\text{II.66.7})$$

For each $i \notin B$, the polynomial $P_{j,B \cup \{i\}}$ is affine in y^j . Its coefficient of y_i^j is $h_B b_j$. Fix an index j_0 with $b_{j_0} \neq 0$. In the equations

$$\sum_j P_{j,B \cup \{i\}}(y^j) = 0, \quad i \notin B,$$

the columns indexed by $y_l^{j_0}$, $l \notin B$, form a diagonal $t \times t$ matrix with nonzero diagonal $h_B b_{j_0}$. Here all other variables are held fixed, including the coordinates whose indices lie in B . Thus these equations have rank at least t and admit at most p^{qk-t} choices of y , for every fixed r .

Suppose first that b is not proportional to s . If $\sum_j b_j \neq 0$, there are no successful arrangements. Otherwise $b \cdot r = 0$ is an independent linear restriction on $H = \ker(s \cdot)$. There are p^{q-2} possible cross-sections, and the preceding rank bound gives probability at most $p^{-1}p^{-t}$.

Now suppose $b = cs$ with $c \neq 0$. Pair each of the preceding affine equations with its vertical counterpart

$$\sum_j r_j P_{j,B \cup \{i\}}(y^j) = 0.$$

If $r_a \neq r_b$, the two columns y_i^a, y_i^b in this pair form the block

$$h_{BC} \begin{pmatrix} s_a & s_b \\ r_a s_a & r_b s_b \end{pmatrix},$$

whose determinant is nonzero. Different indices $i \notin B$ give diagonal blocks, since a variable y_l^j with $l \notin B$ occurs in $P_{j, B \cup \{i\}}$ only when $l = i$. The paired equations therefore have rank at least $2t$ whenever r is not constant. For constant r the rank is still at least t . There are p constant cross-sections among p^{q-1} members of H , so averaging gives the second expression in (II.66.4).

Finally, the first expression is at most p^{-2} , since $t \geq 1$. The second decreases as t increases; at $t = 1$ it is

$$(1 - p^{-(q-2)})p^{-2} + p^{-(q-1)} = p^{-2} + (p-1)p^{-q}.$$

This is (II.66.5). □

Proposition II.66.3 (Sharpness for an individual certificate). *For every $k \geq 1$, $d \geq 2$, and odd prime p , the bound $Q_d(p)$ is attained by a nonparity certificate, conditional on any fixed nonzero sidelengths.*

Proof. Fix a coordinate i and take

$$\eta(\epsilon, j) = s_j \pi(\epsilon) \mathbf{1}_{\{\epsilon_i=0\}}.$$

Parity summation in the remaining $k-1$ coordinates makes every test automatic except

$$\sum_j s_j y_i^j = 0, \quad \sum_j s_j r_j y_i^j = 0.$$

The tests without y_i^j reduce to $\sum_j s_j = 0$ and $s \cdot r = 0$. The displayed pair has rank two for nonconstant r and rank one for constant r . Counting as in Lemma II.66.2 gives exactly $Q_d(p)$. □

[write] **6 October 2026; sources 21–27.** Re-derived independently by source 25 (Theorem II.88.1) and source 27 (Proposition II.94.3, every characteristic), and reproved with credit by source 22 (Lemma II.82.3, degree one).

Remark II.66.4. The restriction $d \geq 2$ is essential for the saving $O(p^{-2})$. If $d = 1$, the constraint $r_1 - r_2 = 0$ forces every cross-section to be constant. The certificate in Proposition II.66.3 then succeeds with probability p^{-1} .

II.66.3 Counting different certificate classes

A direct union bound over all 3^M certificates counts many identical tests. Their most immediate equivalences can be removed exactly.

Lemma II.66.5 (Affine classes of ternary patterns). *Let $p \geq 5$, and let $e \in \{-1, 1\}^n$. Among ternary vectors outside $\text{span}\{e\}$, identify u and u' when $u' = au + be$ for some $a \in \mathbb{F}_p^\times$ and $b \in \mathbb{F}_p$. The number of equivalence classes is*

$$A_n = \frac{3^n - 2^{n+1} + 1}{2}. \tag{II.66.8}$$

Proof. Dividing coordinatewise by e reduces the equivalence to $v' = av + b\mathbf{1}$. Constant patterns are excluded. A pattern taking exactly two values determines an unordered nontrivial partition of the n coordinates. Conversely any two patterns with that partition are affinely equivalent. These give $(2^n - 2)/2$ classes.

For a pattern using all three values, an affine equivalence must permute $\{-1, 0, 1\}$. Summing these three values before and after the affine map gives $3b = 0$, so $b = 0$ because $p \geq 5$. It follows that $a = 1$ or $a = -1$. There are $3^n - 3 \cdot 2^n + 3$ patterns using all three values, and reflection pairs them without fixed points. Adding their classes to the two-value classes proves (II.66.8). $\square$

Since η_0 is always successful, certificates related by

$$\eta' = a\eta + b\eta_0, \quad a \neq 0,$$

have precisely the same success locus. The pure-parity subspace is preserved by this equivalence. After division by η_0 , its members are patterns constant on each of the q cube slices. There are therefore A_q nontrivial pure-parity classes, and exactly

$$J_{k,d} = A_{q2^k} - A_q \quad (\text{II.66.9})$$

nonparity classes. This count may still include impossible certificates and further coincident success loci; it is a safe coefficient for a union bound.

II.66.4 The hyperplanes that determine the leading term

Lemma II.66.6 (Pure-parity hyperplanes). *Let $p > 2d$ be prime. For nonzero sidelengths, a pure-parity certificate $\eta(\epsilon, j) = c_j\pi(\epsilon)$ is successful exactly when*

$$\sum_j c_j = 0, \quad c \cdot r = 0.$$

After excluding $c \in \text{span}\{s\}$, the number of distinct hyperplanes $c \cdot r = 0$ in H obtained this way is

$$C_d = \frac{T_{2d} - 2\binom{2d}{d} + 1}{2}, \quad T_{2d} = \sum_{a=0}^d \binom{2d}{2a} \binom{2a}{a}. \quad (\text{II.66.10})$$

Proof. Parity kills every proper horizontal monomial. Summing the full horizontal monomial gives $(-1)^k h_1 \cdots h_k \sum_j c_j$; adding the last coordinate gives the same nonzero prefactor times $\sum_j c_j r_j$. This proves the characterization.

Write $q = 2d$. Because $p > q$, the equality $\sum_j c_j = 0$ for $c_j \in \{-1, 0, 1\}$ holds in the field if and only if it holds in the integers. There are T_q such vectors: choose the $2a$ nonzero coordinates and then choose which a are positive.

Set $v_j = s_j c_j$. The condition is now $\sum_j s_j v_j = 0$, and it is preserved by $v \mapsto av + b\mathbf{1}$. For a pattern taking two distinct values, let I be the set of coordinates taking one of them. Admissibility is equivalent to $\sum_{j \in I} s_j = 0$. There are

$$\sum_{a=0}^d \binom{d}{a}^2 = \binom{2d}{d}$$

such subsets, including the empty and full sets. Thus there are $\binom{2d}{d} - 2$ two-value classes. The number of admissible ternary patterns using all three values is

$$T_{2d} - 3 \binom{2d}{d} + 3.$$

As in Lemma II.66.5, their classes have size two. Adding the two contributions proves (II.66.10).

It remains to verify that this is a count of distinct hyperplanes, not just of certificate classes. Two nonzero linear functionals on H have the same kernel exactly when they are proportional. The annihilator of H in the ambient dual space is $\text{span}\{s\}$. Hence c and c' define the same hyperplane on H exactly when $c' = ac + bs$ with $a \neq 0$, the equivalence already counted. $\square$

For reference, the first few constants are

d	T_{2d}	$\binom{2d}{d}$	C_d
2	19	6	4
3	141	20	51
4	1107	70	484
8	5 196 627	12 870	2 585 444

Proposition II.66.7 (Growth of the leading coefficient). *Set $C_0 = C_1 = 0$. The ordinary generating function of the constants in (II.66.10) is*

$$\sum_{d \geq 0} C_d z^d = \frac{1}{4} \left((1 - 2\sqrt{z} - 3z)^{-1/2} + (1 + 2\sqrt{z} - 3z)^{-1/2} \right) - (1 - 4z)^{-1/2} + \frac{1}{2(1 - z)}.$$

The paired square-root expressions have only integer powers of z . Moreover,

$$C_d \sim \frac{\sqrt{3}}{4\sqrt{2\pi d}} 9^d \quad (d \rightarrow \infty).$$

Proof. The formula for T_n in (II.66.10) gives

$$\sum_{n \geq 0} T_n t^n = \frac{1}{1 - t} \sum_{a \geq 0} \binom{2a}{a} \left(\frac{t}{1 - t} \right)^{2a} = (1 - 2t - 3t^2)^{-1/2}.$$

Taking its even part and combining the three terms in $C_d = (T_{2d} - 2\binom{2d}{d} + 1)/2$ proves the generating function. For the asymptotic, the constant-term integral is

$$T_{2d} = \frac{1}{2\pi} \int_{-\pi}^{\pi} (1 + 2\cos \theta)^{2d} d\theta.$$

The absolute value of the base is strictly less than 3 away from $\theta = 0$, and near zero

$$\log \frac{1 + 2\cos \theta}{3} = -\theta^2/3 + O(\theta^4).$$

The Gaussian scaling $\theta = u/\sqrt{d}$ therefore gives $T_{2d} \sim \sqrt{3} 9^d / (2\sqrt{2\pi d})$. The central binomial term is $O(4^d)$ and is exponentially smaller. Dividing by two proves the stated equivalent. $\square$

II.66.5 Finite bounds and the optimal leading coefficient

Theorem II.66.8 (Degenerate-arrangement count). *Let $k \geq 1$, $d \geq 2$, and let $p > 2d$ be prime. Write $\delta_{k,d}(p)$ for the proportion of degenerate d -arrangements in $\mathbb{F}_p^{k+1}$, and define z_k , $J_{k,d}$, Q_d , and C_d by (II.66.3), (II.66.9), (II.66.5), and (II.66.10). Then*

$$\delta_{k,d}(p) \leq z_k(p) + (1 - z_k(p)) \left(\frac{C_d}{p} + J_{k,d} Q_d(p) \right) \quad (\text{II.66.11})$$

$$\leq \frac{k + C_d}{p} + J_{k,d}(p^{-2} + (p - 1)p^{-2d}), \quad (\text{II.66.12})$$

whereas

$$\delta_{k,d}(p) \geq \frac{k + C_d}{p} - \frac{\binom{k+C_d}{2}}{p^2}. \quad (\text{II.66.13})$$

Consequently, for fixed k, d and primes $p \rightarrow \infty$,

$$\delta_{k,d}(p) = \frac{k + C_d}{p} + O_{k,d}(p^{-2}). \quad (\text{II.66.14})$$

Equivalently, the number of degenerate arrangements is

$$(k + C_d)p^{(2d+1)k+2d-2} + O_{k,d}(p^{(2d+1)k+2d-3}).$$

Proof. The event of a zero sidelength has probability $z_k(p)$ and consists entirely of degenerate arrangements. Conditional on nonzero sides, Lemma II.66.6 and a union bound give at most C_d/p for pure-parity certificates. Lemma II.66.2 gives at most $Q_d(p)$ for each of the $J_{k,d}$ nonparity classes. This proves (II.66.11). Using $z_k(p) \leq k/p$ and discarding the survival factor proves (II.66.12).

For the lower bound, consider the k zero-side hyperplanes and the C_d pure-parity hyperplanes as subsets of the linear parameter space

$$\mathbb{F}_p^k \times H \times \mathbb{F}_p^{qk}.$$

They are distinct hyperplanes, each with probability $1/p$. Two distinct ones have independent defining linear functionals, so their intersection has probability $1/p^2$. All these hyperplanes consist of degenerate arrangements, including when sidelengths vanish. The first two terms of inclusion-exclusion therefore give (II.66.13). Together, the upper and lower bounds prove (II.66.14), with an explicit error constant if desired. $\square$

[write] **6 October 2026; sources 15–20.** *Against the formal Lemma 15.4* (checked at HEAD 0b4890a1c). lemma_15_4 (Sections14_15.lean, line 402; exact companion lemma_15_4_holds, Proofs15DegenerateArrangements.lean, line 1131) bounds the number of degenerate arrangements by $3^{2d \cdot 2^k} k N^{(2d+1)k+2d-2}$ for odd primes and all $d \geq 1$, with exactly this degeneracy predicate (`GeneralArrangement.IsDegenerate`, line 354: a ternary dependence that is not a modular multiple of the parity dependence and annihilates every monomial). This theorem covers $d \geq 2$ and $p > 2d$ and gives the sharp first-order constant $k + C_d$; for those cases its finite bound (II.66.12) is below the formal one (checked at the write: $J_{k,d} < 3^{2d \cdot 2^k} / 2$, $C_d < 3^{2d} / 2$ and $p \geq 5$ give $(k + C_d)/p + J_{k,d}(p^{-2} + (p - 1)p^{-2d}) < k 3^{2d \cdot 2^k} / p$). It is not formalized. The case $d = 1$, which the remark after Proposition II.66.3 sets aside, is addressed by source 22 (placed by a902613c9, write pending), whose placement records the leading coefficient $(k^2 + 3k)/p$; sources 25 and 27 (placed by 79d9075b9, write pending) re-derive the constant $k + C_d$ independently, and source 27's placement records the second coefficient for $d = 2$. The intake's Monte Carlo of the exact predicate ($k = 1$, $d = 2$) gave $p\delta = 5.43, 5.29, 5.05, 5.14$ at $p = 23, 47, 101, 199$, consistent with $k + C_2 = 5$.

[write] **6 October 2026; sources 21–27.** *Written in the fifth write.* Sources 25 and 27 re-derive the coefficient $k + C_d$ independently, as $k + H_d$ (Theorem F.30.1, over prime-power fields, with the codimension-one geometry) and $k + L_d$ (Theorem F.32.1 and, in every characteristic, Theorem II.96.3); this theorem is the first print. Source 27 determines the second coefficient for $d = 2$, $(k^2 + 21k - 6)/2$ (Theorem F.32.3), and source 22 the leading coefficient $(k^2 + 3k)/p$ for $d = 1$ (Theorem II.85.1) and a finite bound with a degree-one count polynomial in k in place of $J_{k,d}$ (Theorem II.84.1). Section II.81.

Gowers's bound has coefficient $k3^{2d2^k}$ in front of $p^{(2d+1)k+2d-2}$. Theorem II.66.8 identifies the optimal first-order coefficient as $k + C_d$. The extra $J_{k,d}p^{-2}$ term must be retained when using the theorem at a finite prime. Thus one should not simply replace the old coefficient by $k + C_d$ in every finite-size inequality. The next corollary shows how to use the full bound in the intended application.

II.66.6 An explicit replacement for the size condition in Lemma 15.5

For a map ϕ defined on all the vertices of an arrangement, the statement that it *respects* the arrangement means $\sum_{\epsilon,j} \eta_0(\epsilon, j)\phi(\rho(\epsilon, j)) = 0$. Equivalently, the alternating sums of ϕ on the first d cubes add to the corresponding sum on the last d cubes.

Corollary II.66.9 (Refined arrangement selection). *Let $k \geq 1$, let $0 < \alpha, \beta, \eta \leq 1$, and let $p \geq 17$ be prime. Suppose that $B \subseteq \mathbb{F}_p^{k+1}$ has size βp^{k+1} and that $\phi : B \rightarrow \mathbb{F}_p$ respects at least $\alpha\beta^{15}p^{17k+15}$ eight-arrangements in B . Set*

$$M = 2^{k+4}, \quad E_k = 2^{2^{k+4}+k+3}, \quad \lambda = (\alpha\eta/4)^{E_k}\beta^{15},$$

and

$$a_k = k + 2\,585\,444, \quad J_k = A_{2^{k+4}} - A_{16}, \quad A_{16} = 21\,457\,825.$$

If

$$z_k(p) + (1 - z_k(p)) \left(\frac{2\,585\,444}{p} + J_k(p^{-2} + (p-1)p^{-16}) \right) \leq \lambda, \quad (\text{II.66.15})$$

then B has a subset B' containing at least λp^{17k+15} eight-arrangements, of which a proportion at least $1 - \eta$ is respected by ϕ .

An explicit sufficient condition for (II.66.15) is

$$p \geq \frac{a_k + \sqrt{a_k^2 + (72/17)J_k\lambda}}{2\lambda}. \quad (\text{II.66.16})$$

A simpler sufficient condition, using the stated parameter ranges, is

$$\boxed{\lambda p \geq k + 2\,585\,445.} \quad (\text{II.66.17})$$

In particular, (II.66.17) replaces the source condition $\lambda p \geq k3^{2^{k+4}}$ without changing the conclusion or the exponent E_k .

Proof. We first recall the part of Gowers's selection argument that does not use a lower bound on p . For a positive integer R , choose independent uniform field elements t_ℓ and $u_{A,\ell}$ for $1 \leq \ell \leq R$ and $A \subseteq [k+1]$. Given these choices, retain each $x \in B$ independently with probability

$$\prod_{\ell=1}^R \frac{1}{2} \left(1 + \cos \frac{2\pi}{p} \left(t_\ell \phi(x) + \sum_{A \subseteq [k+1]} u_{A,\ell} \prod_{i \in A} x_i \right) \right).$$

The field element inside the cosine may be represented by any integer residue. Expanding one factor as $(2 + \omega^v + \omega^{-v})/4$ and averaging the phase parameters shows that an unrespected nondegenerate arrangement is retained with probability 2^{-MR} . A respected arrangement is retained with probability at least

$$2^{-MR}(1 + 2^{1-M})^R.$$

For a nondegenerate arrangement the only surviving ternary certificates are $0, \eta_0, -\eta_0$; its vertices are distinct, since a repeated pair would give a two-point certificate. For a respected degenerate arrangement these same terms still contribute, and all further surviving terms have nonnegative coefficients, so the inequality remains valid. If vertices repeat, the actual joint retention probability is at least the product with vertex multiplicities, since every individual retention factor lies in $[0, 1]$.

Let X count respected arrangements in the selected set, and let Y count its unrespected nondegenerate arrangements. There are at most $\beta^{15}p^{17k+15}$ eight-arrangements in B : choose the first fifteen base vertices in B , the sidelengths in $\mathbb{F}_p^k$, and the last horizontal base point in $\mathbb{F}_p^k$; the last cross-section is forced. It follows that

$$\eta\mathbb{E}X - \mathbb{E}Y \geq 2^{-MR}\beta^{15}p^{17k+15}(\alpha\eta(1 + 2^{1-M})^R - 1).$$

Choose

$$R = \left\lceil 2^{M-1} \log_2 \frac{2}{\alpha\eta} \right\rceil.$$

The inequality $1 + x \geq 2^x$ for $0 \leq x \leq 1$ gives $\alpha\eta(1 + 2^{1-M})^R \geq 2$. Since $E_k = M2^{M-1}$, the choice of R also gives

$$2^{-MR} \geq 2^{-M}(\alpha\eta/2)^{E_k} \geq 2(\alpha\eta/4)^{E_k};$$

the final inequality is equivalent to $E_k \geq M + 1$, which holds here. With $L = \lambda p^{17k+15}$, some realization therefore satisfies $\eta X - Y \geq 2L$.

By Theorem II.66.8 and (II.66.15), the total number of degenerate arrangements in the entire ambient space is at most L . Thus the number Z of unrespected degenerate arrangements in this realization is at most L . We obtain

$$Y + Z \leq \eta X - L \leq \eta X, \quad X \geq L.$$

The total number of arrangements is at least X , and its bad proportion is at most $\eta X/(X+Y+Z) \leq \eta$. This proves the conclusion.

For the quadratic criterion, $p \geq 17$ gives

$$p^{-2} + (p-1)p^{-16} \leq p^{-2} + p^{-3} \leq \frac{18}{17}p^{-2}.$$

The left side of (II.66.15) is consequently at most $a_k/p + (18/17)J_k/p^2$. Solving the resulting quadratic inequality gives (II.66.16).

Finally, the parameter ranges imply $\lambda \leq 4^{-E_k}$. Since $J_k < 3^M/2$ and $E_k \geq M$,

$$J_k\lambda \leq \frac{1}{2}(3/4)^M \leq \frac{1}{2}.$$

If $\lambda p \geq a_k + 1$, then

$$\frac{a_k}{p} + \frac{18J_k}{17p^2} \leq \lambda \left(\frac{a_k}{a_k + 1} + \frac{(18/17)J_k\lambda}{(a_k + 1)^2} \right) \leq \lambda.$$

This proves (II.66.17). □

[write] **6 October 2026; sources 15–20.** *Against the formal Lemma 15.5* (HEAD 0b4890a1c). `lemma_15_5` (Sections14_15.lean, line 414) is an $\exists N_0$ statement with the same exponent (`arrangementSelectionExponent`, $2^{2^{k+4}+k+3}$) and conclusion. Its proof `lemma_15_5_holds` (Proofs15Restriction.lean, line 3086) takes N_0 from $C/(\delta \cdot \text{scale})$ with $C = 3^{16 \cdot 2^k} k$ in the private `selectionDegenerate_real_small` (line 2690), which is the source condition $\lambda p \geq k 3^{2^{k+4}}$; this corollary would sharpen that internal constant to $k + 2,585,445$. The formal statement does not display N_0 , so its statement does not change; not formalized. Source 22 (placed, write pending) gains nothing over this corollary at leading order at Lemma 15.5, as its placement records.

[write] **6 October 2026; sources 21–27.** Source 22’s thresholds have the same leading term $(k + C_8)/\lambda$ or a worse one (note after its Corollary II.84.5). Source 25’s Theorem II.92.1 extends the selection to every $d \geq 2$ and every finite field of characteristic $p > 2d$, keeping the exact number of Riesz factors.

Remark II.66.10 (What has and has not propagated). Corollary II.66.9 is a proved improvement of one ambient-size requirement in the density-increasing selection. It retains the same Riesz-product exponent, the same guaranteed number of arrangements, and the same error proportion. A final improved quantitative Szemerédi theorem would require propagating all later requirements and verifying which one dominates. That further conclusion is not asserted here.

II.66.7 Further questions suggested by the proof

The next term in (II.66.14) is a concrete finite problem. Intersections of the $k + C_d$ leading hyperplanes contribute at order p^{-2} , and Proposition II.66.3 shows that nonparity certificates also contribute on that scale. Computing the exact coefficient requires accounting for coincident and intersecting success loci, rather than merely counting certificate classes.

A second direction is to preserve the full parameter $t = k - m$ in Lemma II.66.2. A certificate whose first nonzero moment has low degree imposes considerably more than two conditions. Counting certificates according to this degree and according to their moment vectors could reduce $J_{k,d}$ substantially in finite bounds.

The pure-parity hyperplanes form a finite arrangement of balanced subset-sum equations on H . Determining its intersection lattice would give an exact counting polynomial for the corresponding nondegenerate cross-sections away from exceptional characteristics. The proof of Lemma II.66.6 already supplies its number of distinct hyperplanes.

For formalization, the natural first units are the Boolean moment expansion, the diagonal minors in Lemma II.66.2, and the finite-field counts for kernels and constant cross-sections. The affine-pattern enumeration and the two-term inclusion-exclusion bound can then be developed separately. No Lean verification of these new statements is claimed in this article.

II.67 The finite common-sampling problem

Source 19, delivered Section 2.

II.67.1 Rows, classes, and the recovery rule

Let J be a finite set with $n = |J| \geq 1$, and let H be a nonempty finite set of rows. For each $h \in H$, let

$$C_1(h), \dots, C_q(h) \subseteq J$$

be pairwise disjoint; empty classes are allowed. Their union may be a proper subset of J . This last point is essential in applications, because the original function is defined only on a subset of the ambient box. All loss fractions below use the ambient normalization $|H|n$, not the possibly much smaller size of that partial domain.

Fix an integer $d \geq 1$. A subset $R \subseteq J$ is the common set of sampled columns, also called the *anchor set*. A class C is *recovered* when

$$|R \cap C| \geq \min(d, |C|). \quad (\text{II.67.1})$$

If $|C| \geq d$, this provides d distinct interpolation nodes. If $0 < |C| < d$, it says that every point of C is sampled. In that case the interpolant of degree at most $|C| - 1$ agrees with the function on the entire class, which is all that the covering argument needs.

Define the normalized loss

$$L_H(R) = \frac{1}{|H|n} \sum_{h \in H} \sum_{t=1}^q |C_t(h)| \mathbf{1}_{\{|R \cap C_t(h)| < \min(d, |C_t(h)|)\}}. \quad (\text{II.67.2})$$

Classes are discarded as whole classes. Consequently, every retained class still contains all of its chosen interpolation nodes.

Remark II.67.1 (What the extremal model measures). The loss in (II.67.2) measures this specific recovery certificate. A function might happen to be covered by some other polynomial even when its assigned class fails (II.67.1). Our lower bounds do not rule out methods that exploit such additional coincidences, use different anchors in different rows, or infer structure without this interpolation certificate.

II.67.2 Exact finite expectation

Use the convention $\binom{a}{b} = 0$ if $b < 0$ or $b > a$, and let a sum over an empty index set equal zero. For $0 \leq r \leq n$ and $0 \leq s \leq n$, write

$$P_{n,r,d}(s) = \frac{1}{\binom{n}{r}} \sum_{j=0}^{\min(d,s)-1} \binom{s}{j} \binom{n-s}{r-j}, \quad W_{n,r,d}(s) = \frac{s}{n} P_{n,r,d}(s). \quad (\text{II.67.3})$$

Lemma II.67.2 (Exact expected loss). *If R is uniformly distributed among the r -element subsets of J , then*

$$\mathbb{E}_R L_H(R) = \frac{1}{|H|} \sum_{h \in H} \sum_{t=1}^q W_{n,r,d}(|C_t(h)|).$$

The statement remains valid with any nonnegative row weights w_h summing to one, after replacing the row average by $\sum_h w_h$.

Proof. For a fixed class of size s , the number of r -element samples having exactly j points in that class is $\binom{s}{j} \binom{n-s}{r-j}$. Division by the total number $\binom{n}{r}$ gives the hypergeometric probability. Sum over the failure values $j < \min(d, s)$, multiply by the class mass s/n , and use linearity of expectation. No independence between different rows or between different classes is required. $\square$

II.67.3 A uniform bound that retains class mass

Lemma II.67.3 (Marked-sample estimate). *Let $X = |R \cap C|$, where $|C| = s$ and R is a uniform r -element subset of an n -point set. If $d - 1 \leq r < n$, then*

$$\frac{s}{n} \Pr(X < d) \leq \frac{d(n-r)}{n(r+1)} + \frac{d-1}{n} = \frac{d}{r+1} - \frac{r-d+1}{n(r+1)} \leq \frac{d}{r+1}. \quad (\text{II.67.4})$$

Proof. First choose an $(r+1)$ -element subset U uniformly, and then mark one of its elements z uniformly. The pair $(U \setminus \{z\}, z)$ has the same law as a uniform r -element sample followed by a uniform element of its complement. Put $Y = |U \cap C|$. The event that the marked point belongs to C and that the unmarked sample contains fewer than d points of C has probability

$$\mathbb{E} \left[\frac{s-X}{n-r} \mathbf{1}_{\{X < d\}} \right] = \mathbb{E} \left[\frac{Y}{r+1} \mathbf{1}_{\{1 \leq Y \leq d\}} \right] \leq \frac{d}{r+1}.$$

Also $\mathbb{E}[X \mathbf{1}_{\{X < d\}}] \leq d-1$. Adding the two contributions to $s \Pr(X < d) = \mathbb{E}[(s-X) \mathbf{1}_{\{X < d\}}] + \mathbb{E}[X \mathbf{1}_{\{X < d\}}]$ proves the first inequality. The displayed identity and the last inequality are elementary algebra, using $r \geq d-1$. $\square$

Theorem II.67.4 (Common anchors with controlled lost mass). *For every $0 \leq r \leq n$, there is an r -element set $R \subseteq J$ such that*

$$L_H(R) \leq \min \left\{ 1, \frac{qd}{r+1} \right\}. \quad (\text{II.67.5})$$

If $r = n$, the loss is zero. In particular, for $0 < \sigma < 1$, the choice

$$r = \min \left\{ n, \left\lceil \frac{qd}{\sigma} \right\rceil - 1 \right\} \quad (\text{II.67.6})$$

guarantees $L_H(R) \leq \sigma$.

Proof. For $d-1 \leq r < n$, the failure event in (II.67.1) is contained in $\{X < d\}$. Apply lemma II.67.3 to each class and sum. The expectation is at most $qd/(r+1)$, so some sample attains that bound. The loss is always at most one because the classes in a row are disjoint. If $r < d-1$, then $qd/(r+1) \geq 1$, so the trivial bound suffices. If $r = n$, every class is recovered.

For (II.67.6), either all columns are chosen, or $r+1 = \lceil qd/\sigma \rceil \geq qd/\sigma$, which gives the stated loss bound. The sample size is a nonnegative integer; indeed it is at least one unless the full population itself is empty, which was excluded. $\square$

[write] 6 October 2026; sources 15–20. *Against the formal Section 16* (checked at HEAD 0b4890a1c). The formal development has since confirmed this source’s audit point (Section V.80.1) that two hits need not be two distinct anchors: `section16_printed_distinct_sampling_counterexample` (Proofs16DistinctSampling.lean, line 204; commit 94ea4bcc3, after this source’s pin and three minutes before its arrival, so independent of it). Its repair `section16_distinct_sampling_budget` (line 117) samples with replacement and bounds the failure probability of each class of at least two points and density at least σ/q by σ once $r\sigma^2 \geq 6q$; this theorem bounds the total lost mass without replacement, with r of order q/σ instead of q/σ^2 . The formal results do not prove this theorem, and nothing of this source is formalized.

[write] 6 October 2026; sources 21–27. Source 23’s Lemma II.103.8 uses the same add-one argument at $d = 2$ for a different loss (sampled points retained); its stated sample size is one larger than the one here (note there).

Corollary II.67.5 (Weighted rows and nonuniform class counts). *If row h has at most q_h classes and weight $w_h \geq 0$, with $\sum_h w_h = 1$, the same argument gives weighted loss at most*

$$\min \left\{ 1, \frac{d}{r+1} \sum_h w_h q_h \right\}.$$

The exact weighted expectation in lemma II.67.2 may be used instead when class sizes are known.

This estimate avoids the need to discard every small class before sampling. The contribution of a class that is difficult to hit is automatically multiplied by its small mass. For the affine step of Lemma 16.10, it changes the sample scale from $q\sigma^{-2}$ to $q\sigma^{-1}$. Sampling without replacement also ensures that different anchors are distinct, which is necessary for secant interpolation over a field.

II.68 An exact minimax theorem and a deterministic algorithm

Source 19, delivered Section 3.

II.68.1 The finite extremal value

Definition II.68.1. For integers $n, q, d \geq 1$ and $0 \leq r \leq n$, define

$$\mathcal{M}_{n,q,d}(r) = \max_{\substack{s_1, \dots, s_q \in \mathbb{Z}_{\geq 0} \\ s_1 + \dots + s_q \leq n}} \sum_{i=1}^q W_{n,r,d}(s_i). \quad (\text{II.68.1})$$

The feasible set is finite, so the maximum is attained. A maximizing vector records the sizes of the classes; it does not prescribe how classes are arranged across different rows.

Theorem II.68.2 (Exact common-anchor minimax value). *Let the supremum below range over every nonempty finite row set H and every family of at most q pairwise disjoint classes in each row. Then*

$$\sup_{H, (C_t(h))} \min_{\substack{R \subseteq J \\ |R|=r}} L_H(R) = \mathcal{M}_{n,q,d}(r). \quad (\text{II.68.2})$$

In fact, a finite family of rows attains the supremum, and in that family every r -element anchor set has exactly the same loss.

Proof. For the upper bound, each row's class sizes form a feasible vector in (II.68.1). The exact expectation in lemma II.67.2 is therefore at most $\mathcal{M}_{n,q,d}(r)$. At least one sample has loss no larger than its expectation.

For the lower bound, take a maximizing vector $(s_1, \dots, s_q)$ and choose disjoint subsets $C_1^0, \dots, C_q^0$ of J having those sizes. Use all permutations π of J as rows, with $C_i(\pi) = \pi(C_i^0)$. Fix any r -element anchor set R . As π ranges uniformly over all permutations, $|R \cap \pi(C_i^0)|$ has the hypergeometric distribution with population size n , class size s_i , and sample size r . Consequently its average contribution to the loss is exactly $W_{n,r,d}(s_i)$. Summing over i gives $L_H(R) = \mathcal{M}_{n,q,d}(r)$ for every R . Since this construction uses only $n!$ rows, the supremum is attained by a finite family. $\square$

[write] **6 October 2026; sources 34–39.** Source 39 re-derives this theorem for its own rule, which keeps sampled points (its Theorem II.154.2), without citing this source, and solves the profile problem in closed form for that rule (Theorem II.155.3). For the rule here the closed form fails in 199 of 2,835 small cases (note there); the classification asked in source 19's Question 1 stays open for this rule (item II.39).

Remark II.68.3 (The importance of partial rows). The optimization uses $\sum_i s_i \leq n$. If every row had to be completely partitioned into exactly q classes, the constraint would be $\sum_i s_i = n$ and the answer would be a different optimization. In particular, forcing a large leftover class can change the leading constant when q is fixed. Partial domains in the Gowers application make the inequality constraint the appropriate one here.

II.68.2 Exact evaluation by integer dynamic programming

All terms in (II.68.1) have the same denominator $n \binom{n}{r}$. Put

$$u_s = s \sum_{j=0}^{\min(d,s)-1} \binom{s}{j} \binom{n-s}{r-j}, \quad 0 \leq s \leq n,$$

and define integers

$$V_0(b) = 0, \quad V_t(b) = \max_{0 \leq s \leq b} \{u_s + V_{t-1}(b-s)\}, \quad 0 \leq b \leq n.$$

Then

$$\mathcal{M}_{n,q,d}(r) = \frac{V_q(n)}{n \binom{n}{r}}. \quad (\text{II.68.3})$$

The proof is induction on the number of allowed classes: choose the size of the last class and allocate the remaining budget to the preceding classes. Storing an optimizing size at each stage also recovers an extremal profile. The number of arithmetic comparisons is $O(qn^2)$, with $O(n)$ working memory if the profiles themselves are not stored. These operation counts do not count the bit length of the binomial coefficients.

Table II.68.1: Exact finite minimax loss for $n = 32$, $q = 3$, $d = 2$. Each profile is a maximizer of the finite optimization. Unused columns are allowed. Decimal entries are rounded; the rational entries are exact.

r	Maximizing sizes	Exact loss	Decimal loss	Bound $\min(1, 6/(r+1))$
2	(10,11,11)	3553/3968	0.895413	1.000000
4	(10,10,10)	68145/115072	0.592194	1.000000
8	(5,5,5)	34155/115072	0.296814	0.666667
12	(3,3,3)	1539/7936	0.193926	0.461538
16	(2,2,2)	141/992	0.142137	0.352941
24	(2,2,2)	165/1984	0.083165	0.240000
32	(0,0,32)	0	0.000000	0.181818

The zero loss in the last row is not an asymptotic statement. Choosing all columns recovers every class, including singleton classes. The general $qd/(r+1)$ estimate is deliberately uniform and need not be sharp at the finite endpoint.

II.68.3 Choosing the anchors deterministically

The existence proof can be turned into a finite deterministic algorithm whenever the class memberships are supplied. Suppose $A \subseteq J$ is the set of anchors already selected, $a = |A| \leq r$, and let $F(A)$

be the expected unnormalized final lost mass after completing A uniformly to an r -element subset. For a class C , put $s = |C|$ and $b_C = |A \cap C|$. Its contribution to $F(A)$ is

$$\frac{s}{\binom{n-a}{r-a}} \sum_{j=0}^{\min(d,s)-b_C-1} \binom{s-b_C}{j} \binom{n-a-s+b_C}{r-a-j}, \quad (\text{II.68.4})$$

interpreted as zero when the upper limit is negative. The full $F(A)$ is the sum of these contributions over all rows and classes.

For $a < r$, the tower property for finite averages gives

$$F(A) = \frac{1}{n-a} \sum_{x \in J \setminus A} F(A \cup \{x\}). \quad (\text{II.68.5})$$

Choose a column minimizing the right-hand summand, append it to A , and repeat. The conditional expectation never increases. After r choices it equals the actual loss, proving a deterministic version of the sampling theorem. Formula (II.68.4) makes the algorithm exact, using rational arithmetic.

This is an algorithm for explicit finite class data. It does not turn the entire inverse theorem into an efficient algorithm: the number of rows, the structural partition procedures, and the polynomial lists may themselves be enormous.

II.69 Sharp asymptotics of the sampling loss

Source 19, delivered Section 4.

The finite optimization has a clean limiting form when the sample occupies a vanishing fraction of the columns. The critical classes then have size of order n/r , and a sample sees a bounded expected number of their points.

II.69.1 The occupancy constant

For an integer $d \geq 1$, set

$$S_{d-1}(\lambda) = \sum_{j=0}^{d-1} \frac{\lambda^j}{j!}, \quad f_d(\lambda) = \lambda e^{-\lambda} S_{d-1}(\lambda), \quad C_d = \max_{\lambda \geq 0} f_d(\lambda).$$

Lemma II.69.1 (Unique maximizing mean). *There is a unique positive solution λ_d of*

$$S_{d-1}(\lambda) = \frac{\lambda^d}{(d-1)!}. \quad (\text{II.69.1})$$

The function f_d increases up to λ_d and decreases afterward. Thus $C_d = f_d(\lambda_d)$. In particular,

$$\lambda_1 = 1, \quad C_1 = e^{-1}, \quad \lambda_2 = \varphi, \quad C_2 = \varphi^3 e^{-\varphi}.$$

Proof. Differentiation gives

$$f'_d(\lambda) = e^{-\lambda} \left(S_{d-1}(\lambda) - \frac{\lambda^d}{(d-1)!} \right).$$

For $\lambda > 0$, the quotient $S_{d-1}(\lambda)/\lambda^d = \sum_{j=0}^{d-1} \lambda^{j-d}/j!$ is strictly decreasing from $+\infty$ to zero. This proves uniqueness and the sign change. Also $f_d(0) = 0$ and $f_d(\lambda) \rightarrow 0$ at infinity, so the critical point is the global maximum. For $d = 2$, the equation is $1 + \lambda = \lambda^2$, and $1 + \varphi = \varphi^2$ gives the formula. $\square$

Table II.69.1: The first occupancy constants. Values are numerical evaluations of the proved unique-root characterization.

d	λ_d	C_d
1	1.000000000000	0.367879441171
2	1.618033988750	0.839962094657
3	2.269530842081	1.371101604900
4	2.945186161157	1.942380938050
5	3.639547126480	2.543534354087
6	4.349047605111	3.168184815689

II.69.2 Uniform control of the hypergeometric tails

Lemma II.69.2 (A finite exponential-moment bound). *If X is hypergeometric with parameters (n, s, r) and $t \geq 0$, then*

$$\mathbb{E}e^{-tX} \leq \left(1 - \frac{s}{n} + \frac{s}{n}e^{-t}\right)^r \leq \exp\left(-\frac{rs}{n}(1 - e^{-t})\right). \quad (\text{II.69.2})$$

Proof. Assign weight e^{-t} to each point in the class and weight one to each other population point. The expectation is the mean product of r distinct weights, namely the r th elementary symmetric polynomial divided by $\binom{n}{r}$. This mean is at most the r th power of the arithmetic mean of the weights.

For completeness, the last assertion follows by averaging two unequal coordinates while preserving their sum. In the elementary symmetric polynomial, terms containing neither coordinate stay fixed, terms containing exactly one depend only on their sum, and terms containing both are a nonnegative multiple of their product. Replacing the two coordinates by their average cannot decrease the polynomial. Equivalently, on the compact set of nonnegative weights with fixed sum a maximum is attained at equal coordinates; for positive weights and $r \geq 2$, any unequal pair would increase it. The cases $r = 0, 1$ are immediate. This proves the first inequality. The second is $1 - u \leq e^{-u}$. $\square$

Writing $\lambda = rs/n$ and taking $t = 1$, Markov's inequality gives

$$\lambda \Pr(X < d) \leq \lambda \exp(d - 1 - (1 - e^{-1})\lambda). \quad (\text{II.69.3})$$

This tends uniformly to zero as $\lambda \rightarrow \infty$. It is the tail estimate needed to pass from pointwise Poisson convergence to a maximum over all class sizes.

Lemma II.69.3 (Uniform bounded-mean Poisson limit). *Suppose $r \rightarrow \infty$ and $r/n \rightarrow 0$. Uniformly over integers s for which $\lambda = rs/n$ lies in a fixed compact subinterval of $(0, \infty)$, the hypergeometric probabilities satisfy*

$$\Pr(X = j) = e^{-\lambda} \frac{\lambda^j}{j!} + o(1) \quad (j \text{ fixed}).$$

Proof. In the stated range, $s \asymp n/r \rightarrow \infty$, $s/n = O(1/r)$, and $r/n \rightarrow 0$. For sufficiently large parameters, $n - s \geq r$, and

$$\Pr(X = 0) = \prod_{i=0}^{r-1} \left(1 - \frac{s}{n-i}\right).$$

The sum of the quantities $s/(n-i)$ is $\lambda(1 + O(r/n))$, while the sum of their squares is $O(1/r)$, uniformly on the compact interval. Taking logarithms therefore gives $\Pr(X = 0) = e^{-\lambda} + o(1)$.

For fixed j , division by the zero-count probability gives

$$\frac{\Pr(X = j)}{\Pr(X = 0)} = \frac{(s)_j (r)_j}{j! (n - s - r + 1)^{\bar{j}}} = \frac{\lambda^j}{j!} + o(1),$$

where $(a)_j$ and $a^{\bar{j}}$ denote falling and rising factorials. This proves the assertion. In particular, no hypothesis $r^2/n \rightarrow 0$ is required. $\square$

II.69.3 The sharp minimax asymptotic

Theorem II.69.4 (Sharp sparse-sampling law). *Fix $d \geq 1$. Along any sequence of integers n, q, r with*

$$r \rightarrow \infty, \quad r/n \rightarrow 0, \quad q/r \rightarrow 0,$$

one has

$$\frac{r}{q} \mathcal{M}_{n,q,d}(r) \rightarrow C_d. \quad (\text{II.69.4})$$

Proof. For a single class size s , put $\lambda = rs/n$. Then

$$rW_{n,r,d}(s) = \lambda \Pr\{X < \min(d, s)\}.$$

If λ is small, this is at most λ . On any compact interval bounded away from zero, one has $s \rightarrow \infty$, so $\min(d, s) = d$ eventually; lemma II.69.3 gives uniform convergence to $f_d(\lambda)$. For large λ , the failure event is contained in $\{X < d\}$, and (II.69.3) controls the tail uniformly. These three observations imply

$$\max_{0 \leq s \leq n} rW_{n,r,d}(s) \rightarrow C_d.$$

Since there are at most q classes, this proves the upper bound in (II.69.4).

For the lower bound choose $s = \lfloor \lambda_d n/r \rfloor$ and set all q class sizes equal to s . The vector is feasible for all sufficiently large parameters, because $qs/n \leq \lambda_d q/r \rightarrow 0$. Also $rs/n \rightarrow \lambda_d$ since $r/n \rightarrow 0$. The single-class limit gives $rW_{n,r,d}(s) \rightarrow C_d$, and hence $\mathcal{M}_{n,q,d}(r) \geq qW_{n,r,d}(s)$ gives the matching lower bound. $\square$

Corollary II.69.5 (Sharp affine sampling constant). *Under the hypotheses of theorem II.69.4, the affine recovery problem satisfies*

$$\mathcal{M}_{n,q,2}(r) = (0.839962094657 \dots + o(1)) \frac{q}{r}.$$

Thus, in a regime where the desired loss $\sigma \rightarrow 0$ and n is sufficiently large relative to the required sample, the sharp common-sample budget for this recovery mechanism is $(C_2 + o(1))q/\sigma$. The unconditional finite budget $\lceil 2q/\sigma \rceil - 1$ is a simpler uniform guarantee.

[write] **6 October 2026; sources 34–39.** The same constant is source 39's Corollary II.156.3, derived independently; source 39 also gives the limit when q/r stays bounded (its Theorem II.157.3).




Figure II.69.1: Left: the limiting one-class loss profiles and their unique maxima. Right: the exact finite minimax value for $n = 32$, $q = 3$, $d = 2$, compared with the proved elementary bound. The finite curve is computed with integer dynamic programming; it is not a random simulation.

II.69.4 What the asymptotic does not say

The constant C_d concerns partial rows and the sample-based recovery rule (II.67.1). It does not replace the finite formula when r is comparable to n , and it does not assert that an arbitrary polynomial-cover construction requires that many anchors. Nor does it improve the structural complexity of a single sampled slice. The next part of the article addresses how to combine the slices once they have been chosen.

II.70 Recovering polynomial pieces from a common set of slices

Source 19, delivered Section 5.

The point of the next theorem is that all interpolants can be constructed *after one common refinement of the sampled slices*. Refining each pair-generated slope and intercept separately, as in the displayed proof of Gowers' Lemma 16.10, repeats work that the interpolation formula does not require. The distinction affects both the number of refinement steps and the final number of graph candidates.

We allow polynomial degrees to differ between coordinates. This also makes clear that the argument is not special to affine functions. For a field F , a finite set $U \subset F^k$, and a vector $\mathbf{b} = (b_1, \dots, b_k)$ of nonnegative integers, write

$$\mathcal{P}_{\mathbf{b}}(U) = \{P|_U : P \in F[X_1, \dots, X_k], \deg_{X_i} P \leq b_i \text{ for all } i\}.$$

This is a vector space of functions, even when different polynomials induce the same function on U . Multiaffine functions correspond to $\mathbf{b} = (1, \dots, 1)$.

II.70.1 The finite interpolation statement

Let $J = \{\xi_1, \dots, \xi_n\} \subset F$ consist of distinct field elements, with $n \geq 1$, let H be finite and nonempty, and let $A \subset H \times J$. Assume $q, d \geq 1$. Suppose $f : A \rightarrow F$. For each $h \in H$, suppose the row domain

$A_h = \{x : (h, x) \in A\}$ has a partition into at most q nonempty classes $C_{h,t}$, and that $f(h, \cdot)$ agrees on each class with a polynomial of degree less than d , where $d \geq 1$. There is no requirement that the row partitions or their polynomials vary regularly with h .

For a set of distinct anchors $R \subset J$, call a row class C *certified* if

$$|R \cap C| \geq \min(d, |C|).$$

If a class has at least d points, this condition gives d interpolation nodes. If it has fewer points, the condition samples the whole class; interpolating those finitely many values is then enough.

Lemma II.70.1 (Interpolation after a common slice cover). *Let $R = \{\zeta_1, \dots, \zeta_r\} \subset J$, and let $U \subset H$. For each $i \in \{1, \dots, r\}$, suppose functions*

$$g_{i,1}, \dots, g_{i,M_i} : U \rightarrow F$$

belong to a vector space V_U of functions on U . Suppose also that $G \subset U$ has the following property: whenever $h \in G$ and $(h, \zeta_i) \in A$, one of these functions takes the value $f(h, \zeta_i)$ at h .

Then the graph of f , on the certified classes in rows G , is contained in a union of at most

$$\sum_{j=1}^{\min(d,r)} e_j(M_1, \dots, M_r) \quad (\text{L1})$$

graphs of functions in $V_U \otimes_F F[X]_{<d}$. Here e_j denotes the j -th elementary symmetric polynomial. In particular, if $M_i \leq M$, the number of candidates is at most

$$\sum_{j=1}^{\min(d,r)} \binom{r}{j} M^j. \quad (\text{L2})$$

If all classes being retained have at least d points, the single term $e_d(M_1, \dots, M_r)$ suffices.

Proof. For a nonempty set $S \subset \{1, \dots, r\}$ of size at most d , and indices $a_i \in \{1, \dots, M_i\}$ for $i \in S$, form

$$P_{S,a}(h, x) = \sum_{i \in S} g_{i,a_i}(h) \prod_{v \in S \setminus \{i\}} \frac{x - \zeta_v}{\zeta_i - \zeta_v}. \quad (\text{L3})$$

All denominators are nonzero because the anchors are distinct. The displayed function has degree at most $|S| - 1 < d$ in x , and each of its coefficients, as a polynomial in x , belongs to V_U . The number of choices with $|S| = j$ is $\sum_{|S|=j} \prod_{i \in S} M_i = e_j(M_1, \dots, M_r)$.

Fix $h \in G$ and a certified class C . If $|C| \geq d$, choose d of its anchors and let S be their index set. For each $i \in S$, choose a covering function whose value at h is $f(h, \zeta_i)$. The polynomial in (L3), with h fixed, agrees at d distinct points with the degree- $< d$ polynomial supplied for C ; their difference therefore vanishes identically. Thus it agrees with f at every point of C .

If $|C| = s < d$, certification says that every point of C is an anchor. Use all of them in S , and choose the matching covering functions. Formula (L3) agrees with f at every point of C by the interpolation identity itself. Its degree is at most $s - 1$. It need not agree outside this finite class with the originally supplied polynomial, and no such agreement is required.

Taking all S and all index choices proves the graph cover and the counts. $\square$

Remark II.70.2. The lemma works over every field, including fields of small characteristic. It uses distinct-node interpolation and not division by factorials. If $V_U = \mathcal{P}_b(U)$, its candidates have degree bounds $(b_1, \dots, b_k, d - 1)$.

II.70.2 A weighted local lifting theorem

A local slice-cover hypothesis can be stated independently of any particular regularity theorem. Let H have a finite nonnegative measure μ , and let the allowed regions $U \subset H$ have a notion of width. Assume that an allowed region can be partitioned into allowed subregions and that the relevant function spaces restrict to subregions. The product measure on $H \times J$ gives each column counting measure, so its total mass is $n\mu(H)$.

For each fixed $x \in J$, assume the following slice-cover property. For every real scale $m \geq 1$, on every allowed region U of width at least m , and for every $0 < \epsilon < 1$, the partial function $h \mapsto f(h, x)$ has:

- a partition of U into allowed subregions of width at least $m^{\mathfrak{c}(\epsilon)}$;
- an exceptional subset of μ -mass at most $\epsilon\mu(U)$;
- on every subregion, a cover of its remaining graph by at most $\mathfrak{q}(\epsilon)$ functions in its prescribed vector space.

Here $0 < \mathfrak{c}(\epsilon) \leq 1$ and $\mathfrak{q}(\epsilon) \geq 1$. The measure is part of the hypothesis: a counting-measure regularity theorem is not being asserted to hold for arbitrary weights automatically.

Theorem II.70.3 (Shared-slice lifting). *Under the rowwise degree- $< d$ hypothesis above, suppose that the slice-cover property holds on an initial region H of width at least $m \geq 1$. Let $0 < \sigma, \tau < 1$, and set*

$$r = \min\{n, \lceil qd/\sigma \rceil - 1\}, \quad a = \mathfrak{c}(\tau/r), \quad M = \mathfrak{q}(\tau/r). \quad (\text{L4})$$

There is a partition of H into regions of width at least

$$m^{a^r} \quad (\text{L5})$$

and an exceptional subset of $H \times J$, of product mass at most

$$(\sigma + \tau)n\mu(H), \quad (\text{L6})$$

such that, on every product of a final region with J , the remaining graph of f is covered by at most

$$\sum_{j=1}^{\min(d,r)} \binom{r}{j} M^j \quad (\text{L7})$$

functions in the corresponding tensor space $V_U \otimes_F F[X]_{< d}$.

Proof. The finite sampling estimate proved in the preceding section, with row weight $\mu(\{h\})$, supplies r distinct anchors such that the union of uncertified classes has product mass at most $\sigma n\mu(H)$. For completeness, the required bound follows from

$$|C| \Pr\{|R \cap C| < \min(d, |C|)\} \leq \frac{dn}{r+1} \quad (d-1 \leq r < n)$$

and summation over the at most q classes per row. If $r = n$, every class is certified. The choice (L4) therefore applies without any lower bound on n .

Fix such an anchor set $R = \{\zeta_1, \dots, \zeta_r\}$. Refine first for the slice at ζ_1 , using slice error τ/r . In every resulting region refine for the slice at ζ_2 , with the same relative error, and continue through all r slices. At each stage, the exceptional masses sum over a partition of the initial region to at most $(\tau/r)\mu(H)$. The union of the slice exceptions consequently has mass at most $\tau\mu(H)$. Its product with J has mass at most $\tau n\mu(H)$.

At the i -th stage the lower bound for width is raised to the power a . The final width is therefore at least m^{a^r} . Every earlier slice cover survives restriction to the final regions. Thus all r sampled slices have covers on one common partition, with at most M candidates for each slice.

Apply Lemma II.70.1 separately on each final region. Outside the slice exceptions it covers every certified class. The uncertified classes and the products of the slice exceptions with J together have mass at most (L6), and the graph count is exactly the upper bound (L7). $\square$

Remark II.70.4 (Different slice costs). The uniform allocation τ/r is not essential. If the selected slice i has controls $\mathbf{c}_i(\epsilon)$ and $\mathbf{q}_i(\epsilon)$, choose positive errors ϵ_i with $\sum_i \epsilon_i \leq \tau$. The same proof gives width exponent

$$\prod_{i=1}^r \mathbf{c}_i(\epsilon_i)$$

and graph count

$$\sum_{j=1}^{\min(d,r)} e_j(\mathbf{q}_1(\epsilon_1), \dots, \mathbf{q}_r(\epsilon_r)).$$

This retains information about heterogeneous slice complexity. For $\mathcal{P}_{\mathbf{b}}$, the independent degree bounds in $\mathbf{b}$ are preserved as well.

Remark II.70.5 (Why the refinement count is linear in the anchors). The number of possible pairs of anchors is quadratic in r , but the number of slice functions that must be regularized is r . All pair interpolants are fixed linear combinations of functions already present on the common partition. The quadratic quantity belongs in the final candidate count, not in the number of successive partition refinements.

II.70.3 Common differences of arithmetic-progression boxes

The theorem gives Cartesian products of horizontal regions with the column set. Gowers uses boxes whose coordinate progressions all have one common difference. A final synchronization step supplies this geometry with a single explicit factor, instead of paying a geometric loss at every slice refinement.

We use integer arithmetic progressions embedded injectively in the field and subprogressions in their inherited integer coordinates. Thus, if a parent progression has common difference d_0 , a subprogression used in a refinement has step td_0 with a positive integer t . The containment and length inequalities below are integer inequalities before reduction in the field. This is the no-wrap convention needed for the lemma; an arbitrary progression description in a cyclic group without this convention does not supply those inequalities.

Lemma II.70.6 (One final synchronization). *Let $k \geq 1$, and let $P = T \times J$ be a $(k+1)$ -dimensional box of width m , where $m \geq 1$ is an integer and all coordinate progressions have the same inherited integer step. Partition T , if necessary, into boxes whose coordinate lengths belong to $\{m, m+1, \dots, 2m-1\}$. Suppose each of these horizontal boxes has been refined into horizontal boxes of width at least L , where $1 \leq L \leq m$, by subprogressions in the inherited coordinates.*

Then the corresponding Cartesian products with J admit a partition into synchronized $(k+1)$ -boxes of width at least

$$\max\{1, L/4\}. \quad (\text{L8})$$

Every function cover on a horizontal box times J remains valid on its synchronized subboxes, with no increase in its graph count and no additional exceptional points.

Proof. A progression of length at least m can be partitioned into consecutive blocks whose lengths are between m and $2m - 1$: take blocks of length m , and join a final shorter block to its predecessor. This proves the asserted preliminary balancing.

If $L < 4$, singleton boxes give (L8), so assume $L \geq 4$. Let U be a final horizontal box inside a balanced horizontal parent, and let t be the ratio of its common difference to that of the parent. Since each side of U has at least $\lceil L \rceil$ points and the parent's side lengths are at most $2m - 1$, containment gives

$$t(\lceil L \rceil - 1) \leq 2m - 2. \quad (\text{L9})$$

Partition J , in its inherited coordinates, into its nonempty residue progressions of step t . Since $|J| \geq m$, each of these progressions has length at least $\lfloor m/t \rfloor$. By (L9),

$$\left\lfloor \frac{m}{t} \right\rfloor \geq \left\lfloor \frac{\lceil L \rceil - 1}{2} \right\rfloor \geq \frac{\lceil L \rceil}{4} \geq \frac{L}{4}.$$

Here the middle inequality holds for every integer $\lceil L \rceil \geq 4$. The products of U with the residue progressions have the common difference of U , and all their sides have at least $L/4$ points. The products partition $U \times J$. They merely restrict previously constructed graph covers, so neither candidate counts nor exceptional sets grow. $\square$

Corollary II.70.7 (Shared-slice lifting on synchronized boxes). *Suppose the slice-cover hypothesis holds for horizontal arithmetic boxes and their inherited-coordinate refinements. On an initial synchronized box $T \times J$ of width m , the conclusion of Theorem II.70.3 holds with synchronized final boxes of width at least*

$$\max\{1, \tfrac{1}{4}m^{a^r}\}, \quad (\text{L10})$$

the same exceptional proportion, and the same candidate count.

Proof. Balance T into horizontal boxes with side lengths in $[m, 2m - 1]$, and choose the anchor set by averaging over all rows of the original T , with their original weights. Refine the same sampled slices in each balanced box. The proof of Theorem II.70.3 gives horizontal width $L = m^{a^r}$, with errors summing over the balanced partition. Apply Lemma II.70.6 only after all r horizontal refinements are finished. $\square$

II.70.4 An explicit replacement for the local step in Lemma 16.10

We now substitute the parameters used in the pinned source into the theorem. This is a local quantitative statement; it does not presume that its output fits a separately prescribed inductive control function.

At the start of Gowers' Lemma 16.10, put

$$s_0 = \gamma^{-2} s(2^{-(k+2)} \theta, \gamma, k). \quad (\text{L11})$$

The cross-section obtained by fixing the final variable is (γ, s_0) -multiply k -linear. Its horizontal slice controls have the numerical bounds

$$\mathfrak{c}(\epsilon) = c(\epsilon/s_0, \gamma, k)^{s_0}, \quad \mathfrak{q}(\epsilon) = q(\epsilon/s_0, \gamma, k)^{s_0}. \quad (\text{L12})$$

For the geometric conclusion below, assume in addition that these controls are realized by the slice-cover property at every scale using inherited integer-coordinate refinements. This is an explicit extra geometric hypothesis, not a consequence asserted here of the live Lean predicate for proper modular boxes.

Let $S = T \times J$ be a nonempty proper inherited-coordinate box from the preceding line-cover stage, with $k \geq 1$ and width at least $\ell \geq 1$. Work over a field, or over a prime modulus. Put $A = B_1 \cap E \cap S$, where E is the global good set supplied by Lemma 16.9, and let the integer $q_0 \geq 1$ bound the affine graph count on each row of this restricted domain. This is the quantity denoted q_γ , and then q , in the pinned proof. Partition each row domain by assigning a point to the first affine graph containing it; this gives the disjoint classes required above.

Corollary II.70.8 (Local affine lift with shared refinement). *Let $0 < \sigma, \tau < 1$, set $n = |J|$, and define*

$$\begin{aligned} r &= \min\{n, \lceil 2q_0/\sigma \rceil - 1\}, \\ a &= c(\tau/(rs_0), \gamma, k)^{s_0}, \\ M &= q(\tau/(rs_0), \gamma, k)^{s_0}. \end{aligned} \quad (\text{L13})$$

Under the inherited-coordinate slice-cover hypothesis just stated, after deleting at most $(\sigma + \tau)|S|$ additional ambient points, the graph of the function restricted to A has a cover on synchronized boxes of width at least

$$\max \left\{ 1, \frac{1}{4} \ell^{c(\tau/(rs_0), \gamma, k)^{rs_0}} \right\} \quad (\text{L14})$$

by at most

$$r q(\tau/(rs_0), \gamma, k)^{s_0} + \binom{r}{2} q(\tau/(rs_0), \gamma, k)^{2s_0} \quad (\text{L15})$$

multiaffine graph candidates on each box.

Proof. Use degree parameter $d = 2$ in Corollary II.70.7, with the controls (L12). The first term in (L15) covers the certified singleton classes; the second term covers classes having at least two points. Every candidate is affine in the final variable and multiaffine in the preceding k variables, hence multiaffine in all $k + 1$ variables. The identity $a^r = c(\tau/(rs_0), \gamma, k)^{rs_0}$ gives (L14). $\square$

[write] **6 October 2026; sources 15–20.** *Formal analogues, weaker* (HEAD 0b4890a1c; written after this source's pin, partly independently). `section16_local_affine_lift` (`Proofs16LocalAffineLift.lean`, line 13), `section16_local_affine_lift_all_scales` (`Proofs16LiftAllScales.lean`, line 46) and `section16_local_affine_lift_sqrt` (`Proofs16LiftWidth.lean`, line 31) are proved: on a proper box, with $6q \leq r\tau^2$, they give at most $rp + r^2p^2$ multilinear candidates on each final box (here $rM + \binom{r}{2}M^2$), a good set of proportion $1 - 2\tau - \varepsilon$ (here loss $\sigma + \tau$) and width $\sqrt{(m/8)^a}/4$ (here $\frac{1}{4}\ell^a$), with the common slice cover `Section16AnchoredOn.common_slice_cover` (`Proofs16CommonSliceCover.lean`, line 50) and `section16_sampled_slice_cover` (`Proofs16Interpolation.lean`, line 107). The formal construction builds the slice partition inside short parents and does not need this corollary's inherited-coordinate hypothesis; this corollary's width needs it (Lemma II.70.6). Neither is `lemma_16_10` (`Section16.lean`, line 851), which is open.

[write] **6 October 2026; sources 21–27.** Source 23’s Theorems II.103.6 and II.103.7 carry out the geometric part of this step without the inherited-coordinate hypothesis, with width $L_m^A \geq (m/3)^A$ in place of the formal $\sqrt{(m/8)^a}/4$, conditionally on the slice covers they assume; this advances this source’s Question 6 (item II.40). Not formalized.

Here a real number M is simply an upper bound for an integer number of slice candidates. Using the polynomial expression in M remains valid and does not require rounding it upwards. The common-difference conclusion uses the additional inherited-coordinate slice-refinement hypothesis. If only rectangular products are required, Theorem II.70.3 supplies that conclusion without the synchronization step.

For comparison, the displayed source proof uses

$$r_{\text{old}} = \lceil q_0 \sigma^{-2} \rceil, \quad p_{\text{old}} = 4r_{\text{old}}^2 s_0$$

and then applies a multiple-multilinearity cover with refinement and graph exponents p_{old} . In (L14), the corresponding refinement count is $r s_0$, of order $q_0 s_0 / \sigma$. In (L15), the exponent on the slice graph bound is only $2s_0$, with the separate combinatorial factor $\binom{r}{2}$. Thus, up to explicit constants and finite rounding, the refinement count changes from

$$O(q_0^2 s_0 \sigma^{-4}) \quad \text{to} \quad O(q_0 s_0 \sigma^{-1}).$$

This comparison retains q_0 as an input parameter. It does not hide its dependence on earlier parameters in Lemma 16.9.

The original exceptional set supplied by Lemma 16.9 is still present. Its bound is global over the preceding partition; it need not have small relative size in each individual box S . Our theorem only bounds the additional loss from the graph on A . For a total error allowance ρ , one may allocate $\rho/3$ to that initial step, $\rho/3$ to sampling, and $\rho/3$ to the shared slice refinement. The separate deletion of small row classes is unnecessary.

Remark II.70.9 (Limits of the source substitution). The local cover (L14)–(L15) is the proved output. To infer the particular conclusion called $(\gamma, 1)$ -multiple multilinearity in the source, one would still have to compare this output with its fixed functions $c(\rho, \gamma, k+1)$ and $q(\rho, \gamma, k+1)$, uniformly in all outer parameters. The pinned repository version explicitly flags the failure of the displayed older closing comparison. The shared-refinement theorem removes a major avoidable loss, but it does not by itself establish that remaining comparison or a new bound for the full Szemerédi theorem.

Remark II.70.10 (Potential formalization interface). The finite interpolation statement has a direct formalization interface: a finite row partition; a finite anchor subset; evaluation of partial row functions; the Lagrange interpolation identity over a field; and an elementary symmetric count of candidate tuples. The local-cover theorem adds finite partitions, restricted graph covers, and a telescoping bound for exceptional masses. It can be stated independently of the source’s large numerical control functions. This makes the proved local output available for later, separately verified closure calculations.

II.71 Quantitative interpretation and worked comparisons

Source 19, delivered Section 8.

II.71.1 Two independent savings in the interpolation stage

The sampling estimate and the shared refinement address different parameters. Let q_0 be the number of affine classes per row and s_0 the structural iteration parameter of one sampled slice. For a fixed sampling loss σ , the source construction uses $r_{\text{old}} = \lceil q_0 \sigma^{-2} \rceil$ sampled coordinates. Its subsequent iteration parameter is $4r_{\text{old}}^2 s_0$. The common-slice construction uses

$$r = \min\{|J|, \lceil 2q_0/\sigma \rceil - 1\}$$

and a shared-refinement parameter rs_0 . The former improvement is in the choice of the anchors; the latter is in how the anchored information is combined.

For illustration, hold $q_0 = 10$ and $\sigma = 1/10$, and suppose $|J| \geq 1000$. The source sample count is 1000, while the finite guarantee here is 199. The structural iteration parameters are then $4,000,000s_0$ and $199s_0$, respectively. This comparison does not require s_0 to be small. The new affine list has at most

$$199M + 19,701M^2$$

candidates, where the slice list bound is evaluated at its actual new error budget. The figure 19,701 is $\binom{199}{2}$. This is a count of candidate functions, not a claim about runtime.

When q_0 and σ themselves arise from previous stages, their dependence must be substituted before making a global comparison. Likewise, the new and old slice-list functions are evaluated at different arguments. The gain in the number of refinement steps is explicit, but multiplying ratios of unrelated local bounds would not establish a valid final density increment.

II.71.2 General polynomial degree and a vector-valued extension

For degree at most D in the final coordinate, take $d = D + 1$. The sample count is of order $q(D + 1)/\sigma$, and for a common slice list size M the final count is

$$\sum_{j=1}^{D+1} \binom{r}{j} M^j$$

with out-of-range terms interpreted as zero. For fixed D this is a polynomial in rM of degree $D + 1$. A direct use of an abstract closure lemma on every interpolation candidate would instead put the number of candidates into the refinement iteration. The shared-slice theorem avoids that additional iteration.

The algebraic statement also extends to functions taking values in an F -vector space V . Define a polynomial with values in V by specifying its coefficient vectors. Lagrange interpolation uses only scalar multiplication and addition of these vectors, so the proof of lemma II.70.1 is unchanged. For a recovered large class, the scalar Vandermonde matrix is invertible, and hence its action on a vector of coefficients is invertible as well. This gives uniqueness of a vector-valued polynomial of degree less than d from d distinct values. For a smaller recovered class, direct interpolation still gives agreement at every class point.

Over a commutative ring, the same argument is valid if every denominator $\zeta_i - \zeta_j$ used in interpolation is a unit. Distinctness alone is insufficient. For example, modulo 8 the two polynomials 0 and $2X$ agree at $X = 0$ and $X = 4$ but disagree at $X = 1$. Thus an unrestricted composite-modulus version cannot use a two-point uniqueness assertion.

II.71.3 The gap in nonlinear partition exponents

For q quadratic phases, the optional Lau specialization gives the sufficient flexible partition condition

$$(44q^2 + 45q + 1)a + (44q^2 + 23q)b < 1.$$

The independently proved necessary condition is

$$(3q + 1)a + 2qb \leq 1.$$

The difference between quadratic and linear dependence on q is substantial. The recurrence transfer improves an exponential dependence inherited from phase-by-phase refinement, but it does not determine the optimal nonlinear frontier.

The obstruction already applies to the single cell containing the left endpoint of the interval. This is why the number of candidate common differences is only of order N/L . A proof that handles only a positive fraction of the interval could evade this endpoint argument by discarding the endpoint, and would be a different partition problem. Complete coverage is therefore a material part of both the sufficient and necessary theorems.

II.72 Source 24's chapter: Section 7 again [write]

Added in the fifth write, 6 October 2026.

The next eight sections are source 24, a manuscript on Gowers's Section 7: dependent random choice (Lemma 7.4), the elementary Balog–Szemerédi–Gowers step (Proposition 7.3) and the Freiman restrictions (Lemma 7.5, Corollary 7.6). *Logically the chapter follows source 03's Sections II.5–II.6 and source 10's Sections II.20–II.21*, and most of it repeats them: its pin (2e0553b0c, 13:44) predates their placement, so the repetition is independent, but in this report it is a second route. It is printed here, at the end of Part II, because no number printed earlier may move. Its introduction is Section F.29; its examples and checks, Lean integration, questions, conclusion and appendices are in Part V (Sections V.96–V.101). Its notation is in the table at the head of this Part.

What is proved, and what is new in this report. Second routes: the cubic certificate $7\delta^3 n / (3\sqrt{6})$ with the same tangent point $7u/6$ (Theorem II.74.1 = source 03's Theorem II.5.2), the scalar optimum $\sqrt{C_3} = 0.95378\dots$ and the two-sample barrier (Theorem II.75.1, with source 03's κ_* and two-witness remark; source 24 adds the exact infimum statement and the range $u \leq 2/z_t$), the fibre extraction with prefactor $|B|/(kQ_k)$ (Theorem II.79.3; source 10's Theorem II.21.1 has $|B|/(k(h_k - 1) + 1)$, never worse), and the modern-input corollary (Corollary II.80.2, identical to source 10's alternative after Corollary II.21.2). Weaker than an earlier source: Proposition 7.3 with $(\kappa/8, 2^{15}\kappa^{-5}, 2^{17}\kappa^{-6})$ (Corollary II.77.2; better than source 03's $(c/8, 2^{19}c^{-9})$, worse than source 10's $(c/4, 2^{10}c^{-5})$ with relative doubling $2^{12}\gamma^{-6}$), and the order- k restriction $2^{-(34k+20)}\gamma^{12k+7}/k$ (Theorem F.29.1; source 10's Corollary II.21.2 has $2^{-(24k+14)}\gamma^{12k+7}/k$, better by 2^{10k+6} , so $2^{-209}\gamma^{103}$ against $2^{-295}\gamma^{103}$ at order eight). New in this report: the extraction for codomains $\mathbb{F}_p^d$ by a random linear map, retaining $|B|/p^{\lceil \log_p Q_k \rceil} \geq |B|/(pQ_k)$ independently of d (Theorem II.79.5; it beats source 10's torsion-budget Theorem II.21.3 on $\mathbb{F}_p^d$ whenever $Q_k \leq p^{d-1}$ and $k \geq 2$, proved in the note after it; source 24's own suite does not test it), the parameterized moment–defect selector (Theorem II.73.1, Corollary II.73.2; a general form of source 03's Corollary II.5.3), the written-out Petridis proof of Plünnecke–Ruzsa (Section II.78; classical) and the threshold optimizations of its own route.

Relation to the formal statements (checked at HEAD 2dfe46f44). lemma_7_4, lemma_7_5 and corollary_7_6 (Sections06_07.lean, lines 167, 183, 197) have exact companions (Proofs07DRC.lean, line 536; Proofs07FreimanRestriction.lean, line 267; Proofs07AdditiveRestriction.lean, line 100); proposition_7_3 (line 135, for $\mathbb{Z}^n$) has the exact companion proposition_7_3_holds (Proofs07BalogSzemerédi.lean, line 996), and its form in every abelian group is already formal, with the paper's constants: proposition_7_3_group_holds (line 854). The statements here imply the formal conclusions with better constants, as sources 03's and 10's do (notes there); none is formalized.

II.73 Notation and the exact moment–defect principle

Source 24, delivered Section 2.

All sets used in counting arguments are finite. Unless explicitly weighted, pairs and tuples are ordered, and diagonal entries are included. Graphs below are symmetric incidence relations that may have loops. Walks may repeat vertices. These conventions make the identities exact; excluding repetitions would introduce avoidable error terms.

For finite subsets of an abelian group, write

$$r_A(d) = |\{(a, b) \in A^2 : a - b = d\}|, \quad E(A) = \sum_d r_A(d)^2.$$

Equivalently, $E(A)$ counts the ordered solutions of $a + b = c + d$ in A^4 . If $m = |A| > 0$, then

$$\sum_d r_A(d) = m^2, \quad r_A(d) \leq m, \quad 0 < E(A) \leq m^3. \quad (\text{II.73.1})$$

For $k \geq 1$, kA denotes the k -fold sumset; $0A = \{0\}$.

II.73.1 Weighted selection with the full codegree profile

Let (I, μ) and (V, ν) be finite probability spaces, and let $A_i \subseteq V$ for each $i \in I$. Define

$$c(i, j) = \nu(A_i \cap A_j).$$

For a subset $K \subseteq I$ and $0 < \beta \leq 1$, set

$$\text{bad}_\beta(K) = \sum_{\substack{i, j \in K \\ c(i, j) < \beta}} \mu(i)\mu(j).$$

This is a weighted pair mass, not a conditional probability within K .

Theorem II.73.1 (Exact moment–defect selector). *Fix an integer $t \geq 1$ and $0 < \eta < 1$. Put*

$$\begin{aligned} M_t &= \sum_{i, j} \mu(i)\mu(j)c(i, j)^t, \\ B_t &= \sum_{c(i, j) < \beta} \mu(i)\mu(j)c(i, j)^t, \\ \Delta &= M_t - \eta^{-1}B_t. \end{aligned} \quad (\text{II.73.2})$$

If $\Delta > 0$, there is a set $K \subseteq I$ such that

$$\mu(K)^2 - \eta^{-1} \text{bad}_\beta(K) \geq \Delta. \quad (\text{II.73.3})$$

Consequently,

$$\mu(K) \geq \sqrt{\Delta}, \quad \text{bad}_\beta(K) < \eta \mu(K)^2.$$

Moreover, K may be chosen as the indices whose sets contain all the points of one ordered t -sample from V .

Proof. Sample $v_1, \dots, v_t$ independently according to ν , and put $K = \{i : v_1, \dots, v_t \in A_i\}$. For every ordered pair (i, j) ,

$$\Pr(i, j \in K) = c(i, j)^t.$$

Expanding the square of $\mu(K)$ and taking expectations gives $\mathbb{E}\mu(K)^2 = M_t$. The identical calculation restricted to bad pairs gives $\mathbb{E} \text{bad}_\beta(K) = B_t$. Therefore

$$\mathbb{E}[\mu(K)^2 - \eta^{-1} \text{bad}_\beta(K)] = \Delta.$$

At least one sample in the finite support has score at least its expectation. The two consequences follow because the bad mass is nonnegative and Δ is positive. $\square$

The point is to retain B_t as a *truncated moment*. The often-used estimates $M_t \geq u^t$ and $B_t \leq \beta^t$, where $u = \mathbb{E}_{i,j} c(i, j)$, discard the fact that the two terms come from the same distribution. The cubic argument in the next section exploits precisely this dependence.

Corollary II.73.2 (Higher-arity and heterogeneous versions). *For an integer $r \geq 1$, put $c(i_1, \dots, i_r) = \nu(A_{i_1} \cap \dots \cap A_{i_r})$. Replace the pair distribution in (II.73.2) by $\mu^{\otimes r}$ and define the bad r -tuple mass analogously. Then $\Delta > 0$ implies that some t -sample core satisfies*

$$\mu(K)^r - \eta^{-1} \text{bad}(K) \geq \Delta, \quad \mu(K) \geq \Delta^{1/r}.$$

The identity remains valid if “bad” is specified by any fixed set of index tuples, including tuple-dependent intersection thresholds.

Proof. Expand $\mu(K)^r$ as a sum over ordered r -tuples. The probability that all their indices survive is $c(i_1, \dots, i_r)^t$. $\square$

II.73.2 The unweighted form

When $|I| = n$ and $|V| = m$ carry uniform measure, a sample core has score

$$|K|^2 - \eta^{-1} b_\beta(K), \quad b_\beta(K) = |\{(i, j) \in K^2 : |A_i \cap A_j| < \beta m\}|.$$

The expectation is $n^2 \Delta$. Thus the normalization introduces neither an extra density factor nor a loss from unordered pairs. For later use, if the total number of incidences is $p m n$, then Cauchy–Schwarz gives

$$\frac{1}{m n^2} \sum_{i,j} |A_i \cap A_j| = \frac{1}{m n^2} \sum_{v \in V} d(v)^2 \geq p^2, \quad d(v) = |\{i : v \in A_i\}|. \quad (\text{II.73.4})$$

There is no need to remove low-degree indices before applying this inequality.

II.74 A three-sample strengthening of Lemma 7.4

Source 24, delivered Section 3.

Theorem II.74.1 (Cubic dense-intersection theorem). *Let $A_1, \dots, A_n \subseteq V$, with $|V| = m > 0$, $n > 0$, and $0 < \delta \leq 1$. Suppose*

$$\sum_{i,j=1}^n |A_i \cap A_j| \geq \delta^2 mn^2. \quad (\text{II.74.1})$$

Then there is a set $K \subseteq \{1, \dots, n\}$ such that

$$|K| \geq \frac{7}{3\sqrt{6}} \delta^3 n \quad (\text{II.74.2})$$

and at least 90% of the ordered pairs in K^2 satisfy

$$|A_i \cap A_j| \geq \frac{\delta^2 m}{2}.$$

More strongly, if $b(K)$ is the number of pairs violating this inequality, one can arrange

$$|K|^2 - 10b(K) \geq \frac{49}{54} \delta^6 n^2. \quad (\text{II.74.3})$$

The set K is the common incidence core of three points of V .

Proof. Set $u = \delta^2$, and consider

$$h(s) = s^3 - 10s^3 \mathbf{1}_{s < u/2} = \begin{cases} -9s^3, & 0 \leq s < u/2, \\ s^3, & s \geq u/2. \end{cases}$$

We claim that, for every $s \geq 0$,

$$h(s) \geq L(s) := \frac{49}{12} u^2 s - \frac{343}{108} u^3. \quad (\text{II.74.4})$$

Put $b = 7u/6$. If $s \geq u/2$, then

$$s^3 - L(s) = (s - b)^2(s + 2b) \geq 0.$$

For $0 \leq s \leq u/2$, the function $-9s^3 - L(s)$ is decreasing. Its value at $u/2$ is $u^3/108 > 0$. This proves (II.74.4), including the threshold where the definition of h jumps upwards. The parameter b need not lie in $[0, 1]$; the factorization is valid for all nonnegative s .

For $c(i, j) = |A_i \cap A_j|/m$, hypothesis (II.74.1) says $\mathbb{E}_{i,j} c(i, j) \geq u$. Since L has positive slope,

$$\mathbb{E}_{i,j} h(c(i, j)) \geq L(u) = \left(\frac{49}{12} - \frac{343}{108} \right) u^3 = \frac{49}{54} u^3.$$

Apply Theorem II.73.1 with $t = 3$, $\eta = 1/10$, and $\beta = u/2$. It yields (II.74.3), from which both conclusions follow. $\square$

[write] **6 October 2026; sources 21–27.** *Second route.* This is source 03's Theorem II.5.2: the same hypothesis, the same supporting line (tangent to s^3 at $7u/6$, value $u^3/108$ at $u/2$, hence $49/54$) and the same conclusion; proved independently (see the note in Section F.29).

Corollary II.74.2 (Individual density hypothesis). *The same conclusion holds if $|A_i| \geq \delta m$ for every i .*

Proof. The total incidence count is at least δmn , so (II.73.4) implies (II.74.1). $\square$

The retained-size ratio relative to the source's lemma is

$$\frac{7\delta^3/(3\sqrt{6})}{\delta^5/\sqrt{2}} = \frac{7}{3\sqrt{3}}\delta^{-2} > 1.$$

For instance, at $\delta = 1/10$ the improvement factor in the stated lower bound is approximately 134.7. This numerical example is only an illustration of an exact inequality, not evidence for it.

If $n = 0$, choose $K = \emptyset$. If $m = 0$ and $0 < \delta \leq 1$, choose all indices: every intersection meets the zero threshold and $7\delta^3/(3\sqrt{6}) \leq 1$. Thus the theorem can be extended to the zero-cardinality cases required by a fully general finite-set API.

II.75 The sharp scalar relaxation and a two-sample barrier

Source 24, delivered Section 4.

We now solve the scalar optimization problem behind the preceding certificate. The distinction between a scalar distribution and a realizable codegree matrix is important: the latter has additional constraints, including positive semidefiniteness.

II.75.1 An exact infimum for the score

For $t \in \{2, 3\}$, let $z_t > 1$ be the unique solution of

$$(t-1)z_t^t - tz_t^{t-1} - 9 = 0. \quad (\text{II.75.1})$$

The derivative of the left side is $t(t-1)z^{t-2}(z-1) > 0$ for $z > 1$, and its value at $z = 1$ is negative. For both these values of t , the solution satisfies $z_t > 2$. Define

$$C_t = \frac{tz_t^{t-1} - 9}{2^t}. \quad (\text{II.75.2})$$

Theorem II.75.1 (Sharp mean-only score). *Fix $t \in \{2, 3\}$ and $0 < u \leq 2/z_t$. Among all probability laws on $[0, 1]$ with mean exactly u , the infimum of*

$$\mathbb{E} \left[Z^t - 10Z^t \mathbf{1}_{Z < u/2} \right] \quad (\text{II.75.3})$$

is $C_t u^t$. The same value is the infimum when the mean is required only to be at least u . For every $u > 0$, without the restriction $u \leq 2/z_t$, the lower bound $C_t u^t$ still holds for laws on $[0, 1]$ of mean at least u .

Proof. Write $\beta = u/2$ and $b = z_t \beta$. Let

$$h_t(s) = s^t - 10s^t \mathbf{1}_{s < \beta}, \quad \ell(s) = tb^{t-1}s - (t-1)b^t.$$

The line ℓ is tangent to s^t at b , and (II.75.1) gives

$$\ell(\beta) = -9\beta^t.$$

For $s \geq \beta$, convexity of s^t gives $h_t(s) \geq \ell(s)$. For $0 \leq s \leq \beta$, the difference $-9s^t - \ell(s)$ is decreasing and vanishes at β , so is nonnegative. This proves the global pointwise minorant, with the value at the threshold interpreted on the lower branch only for the purpose of the comparison. Consequently, if $\mathbb{E}Z \geq u$, then

$$\mathbb{E}h_t(Z) \geq \ell(u) = -9\beta^t + tb^{t-1}(u - \beta) = C_t u^t.$$

For sharpness assume $b \leq 1$. Since $z_t > 2$, we have $\beta < u < b$. For each sufficiently small $\varepsilon > 0$, place mass

$$q_\varepsilon = \frac{u - (\beta - \varepsilon)}{b - (\beta - \varepsilon)}$$

at b , and the remaining mass at $\beta - \varepsilon$. This law has mean exactly u . As $\varepsilon \downarrow 0$, both support points approach points of contact with the minorant, from the correct side of the jump at β . The objective therefore converges to $\ell(u) = C_t u^t$. The lower bound and the matching sequence of laws prove the infimum statements. $\square$

[write] **6 October 2026; sources 21–27.** *Source 03's constants* (checked at the write). For $t = 3$, (II.75.1) is source 03's equation $2z^3 - 3z^2 = 9$ (II.5.5), and $C_3 = (3z^2 - 9)/8$ equals source 03's $\kappa_*^2 = (3z^2 - z^3)/4$ because $z^3 = (3z^2 + 9)/2$; so $\sqrt{C_3} = 0.9537881585\dots$ is source 03's (II.5.6). For $t = 2$, $z_2/2 = (1 + \sqrt{10})/2$ is source 03's h , and $C_2 = (-7 + 2\sqrt{10})/4 = 2h - h^2$ is source 03's two-witness value. Source 03 states the optimality of its scalar argument and the two-witness barrier in prose; the exact infimum statement, for means exactly u and at least u , and the range $u \leq 2/z_t$ are this source's.

II.75.2 The improved cubic constant

For $t = 3$, the root and coefficient are

$$z_3 = 2.3293557629\dots, \quad C_3 = \frac{3z_3^2 - 9}{8} = 0.9097118514\dots$$

Thus Theorem II.74.1 admits the slightly stronger conclusion

$$|K| \geq \sqrt{C_3} \delta^3 n = 0.9537881585\dots \delta^3 n. \quad (\text{II.75.4})$$

The rational certificate $49/54$ is only slightly weaker than C_3 , and avoids algebraic-root infrastructure in a formalization. The coefficient C_3 cannot be increased using only the mean constraint and this fixed scalar score, since Theorem II.75.1 supplies matching scalar laws for all sufficiently small u .

II.75.3 What exactly fails with two samples

For $t = 2$, equation (II.75.1) gives $z_2 = 1 + \sqrt{10}$, and

$$C_2 = \frac{-7 + 2\sqrt{10}}{4} = -0.1688611699\dots < 0. \quad (\text{II.75.5})$$

For $0 < u \leq 2/(1 + \sqrt{10})$, a mean constraint alone therefore cannot force a positive expectation of the two-sample score $|K|^2 - 10b(K)$ at threshold $u/2$. The near-extremizers in the proof have two codegree levels, just below $u/2$ and near $(1 + \sqrt{10})u/2$.

This is a precise obstruction to one proposed proof strategy, not a counterexample to a two-sample combinatorial theorem. A scalar near-extremizer need not be the distribution of entries of a codegree matrix. It is also not an obstruction to changing the score, changing the sampling law, or lowering the threshold. Indeed, the next section uses one sample successfully at a different threshold.

II.76 From a profile certificate to many four-step walks

Source 24, delivered Section 5.

Let A be a finite subset of an abelian group, with $m = |A| > 0$. Let D be a finite symmetric set of possible edge differences, and join $a, b \in A$ when $a - b \in D$. Write

$$N(a) = \{b \in A : a - b \in D\}, \quad c(a, a') = \frac{|N(a) \cap N(a')|}{m}.$$

Apply the uniform version of Theorem II.73.1 to these neighbourhoods. The following lemma records the full profile without rounding it to a power of the original energy.

Lemma II.76.1 (Profile-to-difference-set certificate). *Suppose some core $K \subseteq A$ satisfies*

$$|K| \geq \sqrt{\Delta} m, \quad b_\beta(K) \leq \eta |K|^2, \quad \Delta > 0,$$

where a pair is bad when $c(a, a') < \beta$. Choose $\eta < \lambda < 1/2$, and put $r = 1 - \eta/\lambda$. Then there is a nonempty $S \subseteq K$ such that

$$|S| \geq r|K| \geq r\sqrt{\Delta} m \tag{II.76.1}$$

and

$$\frac{|S - S|}{m} \leq \frac{(|D|/m)^4}{(1 - 2\lambda)\sqrt{\Delta}\beta^2}, \tag{II.76.2}$$

$$\frac{|S - S|}{|S|} \leq \frac{(|D|/m)^4}{(1 - 2\lambda)(1 - \eta/\lambda)\Delta\beta^2}. \tag{II.76.3}$$

Proof. Remove from K every vertex with more than $\lambda|K|$ bad partners in K . Since the sum of these bad degrees is $b_\beta(K)$, at most $(\eta/\lambda)|K|$ vertices are removed. This gives (II.76.1).

Fix $a, a' \in S$. At least $(1 - 2\lambda)|K|$ vertices $v \in K$ form good pairs with both a and a' . For every such v , there are at least βm choices of $b \in N(a) \cap N(v)$ and at least βm choices of $c \in N(v) \cap N(a')$. Thus there are at least

$$R = (1 - 2\lambda)|K|\beta^2 m^2 \tag{II.76.4}$$

walks a, b, v, c, a' . Each walk determines a tuple

$$(a - b, b - v, v - c, c - a') \in D^4$$

whose sum is $a - a'$. For fixed endpoints, this map is injective: the first three entries successively recover b, v, c .

For each $d \in S - S$, select one endpoint pair with difference d . It contributes at least R distinct elements of the sum fibre $\{(d_1, d_2, d_3, d_4) \in D^4 : d_1 + d_2 + d_3 + d_4 = d\}$. Fibres for different d are disjoint. Hence

$$|S - S|R \leq |D|^4.$$

Use $|K| \geq \sqrt{\Delta} m$ to obtain (II.76.2), and also $|S| \geq r|K|$ to obtain (II.76.3). $\square$

Remark II.76.2 (Why repetitions do not cause a loss). Nothing in the injection requires the vertices of a walk to be distinct. Likewise, the pair $a = a'$ is legitimate. All estimates are therefore exact for small sets as well as asymptotically valid for large ones.

II.76.1 The one-sample specialization

A one-neighbourhood score followed by bad-degree pruning also appears, with different parameters, in Reiher–Schoen, Lemma 2.1 [12]. This selection component is an adaptation of established machinery, not a claim of a new sampling method.

Suppose the graph has ordered edge density

$$p = \frac{1}{m^2} \sum_{a \in A} |N(a)| > 0.$$

By (II.73.4), the mean normalized codegree is at least p^2 . Use

$$t = 1, \quad \beta = p^2/16, \quad \eta = 1/8. \quad (\text{II.76.5})$$

Since $B_1 \leq \beta$, its profile score satisfies

$$\Delta = M_1 - 8B_1 \geq p^2 - 8\beta = p^2/2.$$

Choose $\lambda = 1/4$ in Lemma II.76.1. We obtain

$$|K| \geq pm/\sqrt{2}, \quad |S| \geq |K|/2 \geq pm/(2\sqrt{2}). \quad (\text{II.76.6})$$

For each fixed endpoint pair in S , the number of walks is at least

$$\frac{|K|}{2} \left(\frac{p^2 m}{16} \right)^2 = \frac{|K| p^4 m^2}{512}. \quad (\text{II.76.7})$$

The weaker codegree threshold makes one sample sufficient; there is no contradiction with the two-sample obstruction at threshold $p^2/2$.

II.77 An explicit energy-to-small-difference-set theorem

Source 24, delivered Section 6.

The key improvement in the energy reduction is not merely a more favourable choice of constants in dependent random choice. It is the retention of the *same actual popular-edge mass* in two places: the incidence density and the size of the popular difference set. The resulting powers cancel in the four-walk count.

Theorem II.77.1 (Popular-mass form). *Let A be a nonempty finite subset of an abelian group, $m = |A|$, and $\kappa = E(A)/m^3$. Choose $0 < \tau < \kappa$, and define*

$$D = \{d : r_A(d) \geq \tau m\}, \quad p = \frac{1}{m^2} \sum_{d \in D} r_A(d). \quad (\text{II.77.1})$$

Then

$$p \geq \frac{\kappa - \tau}{1 - \tau}, \quad |D| \leq \frac{p}{\tau} m. \quad (\text{II.77.2})$$

There is a nonempty $S \subseteq A$ with

$$|S| \geq \frac{p}{2\sqrt{2}}m, \quad (\text{II.77.3})$$

$$|S - S| \leq \frac{512\sqrt{2}}{p\tau^4}m, \quad (\text{II.77.4})$$

$$|S - S| \leq \frac{2048}{p^2\tau^4}|S|. \quad (\text{II.77.5})$$

Proof. The difference function is symmetric, so $D = -D$. On D we have $r_A(d)^2 \leq mr_A(d)$, and outside D we have $r_A(d)^2 \leq \tau mr_A(d)$. Summing gives

$$\kappa m^3 \leq m \sum_{d \in D} r_A(d) + \tau m \sum_{d \notin D} r_A(d) = (p + \tau(1 - p))m^3.$$

This proves the first estimate in (II.77.2). The second follows from $\tau m|D| \leq \sum_{d \in D} r_A(d) = pm^2$.

Construct the popular-difference graph. Its ordered edge density is exactly p , not just a lower bound in terms of κ . Apply the one-sample construction of Section II.76, obtaining K, S as in (II.76.6). The four-walk injection gives

$$|S - S| \leq \frac{512|D|^4}{|K|p^4m^2} \leq \frac{512m^2}{|K|\tau^4} \leq \frac{512\sqrt{2}}{p\tau^4}m.$$

For the relative estimate, divide the middle expression by $|S| \geq |K|/2$, rather than by a separately rounded lower bound:

$$\frac{|S - S|}{|S|} \leq \frac{1024m^2}{|K|^2\tau^4} \leq \frac{2048}{p^2\tau^4}.$$

This proves all assertions. $\square$

Corollary II.77.2 (Dyadic constants). *If $E(A) \geq \kappa m^3$ with $0 < \kappa \leq 1$, there is a nonempty $S \subseteq A$ satisfying, simultaneously,*

$$\boxed{|S| \geq 2^{-3}\kappa m, \quad |S - S| \leq 2^{15}\kappa^{-5}m, \quad |S - S| \leq 2^{17}\kappa^{-6}|S|.} \quad (\text{II.77.6})$$

Proof. For equality in the energy hypothesis, set $\tau = \kappa/2$. Then $p \geq \kappa/(2 - \kappa) \geq \kappa/2$. The first two bounds follow from

$$\frac{p}{2\sqrt{2}} \geq \frac{\kappa}{4\sqrt{2}} \geq 2^{-3}\kappa, \quad \frac{512\sqrt{2}}{p\tau^4} \leq 2^{14}\sqrt{2}\kappa^{-5} \leq 2^{15}\kappa^{-5}.$$

The relative estimate is

$$\frac{2048}{p^2\tau^4} \leq 2^{11} \cdot 2^2 \cdot 2^4\kappa^{-6} = 2^{17}\kappa^{-6}.$$

If the assumed energy parameter is smaller than the actual normalized energy, apply the equality case to the latter and then weaken each bound using its monotonicity. $\square$

[write] **6 October 2026; sources 21–27.** *Against sources 03 and 10.* Source 03's Theorem II.6.3 gives $(c/8, 2^{19}c^{-9})$: the same retained size, a weaker difference bound. Source 10's Theorem II.20.1 gives $(c/4, 2^{10}c^{-5})$, hence relative doubling $2^{12}c^{-6}$ (the proof of its Corollary II.21.2): better in all three conclusions. Reiher–Schoen's theorem is stronger than all three, as the source says.

Remark II.77.3 (A constant that must not be recomputed incorrectly). Dividing the *rounded* absolute difference bound by the *rounded* retained-size bound gives only $2^{18}\kappa^{-6}$. The better factor 2^{17} comes from the direct relative estimate in Theorem II.77.1. The main exponent ledger uses that direct estimate.

II.77.1 Why retaining p matters

The coarse bound $|D| \leq m/\tau$ is true, but replacing the second estimate in (II.77.2) by it loses the factor p^4 in $|D|^4$. That factor exactly cancels the p^4 coming from two common-neighbour choices in (II.76.7). Keeping only an energy lower bound for the edge count while discarding the same count in the popular-set estimate would therefore produce a substantially weaker power of κ .

Similarly, pruning the original graph to force a minimum degree is unnecessary here. The mean codegree estimate follows directly from all its incidences, and the only pruning occurs after a sample has produced a core with few bad pairs. This is what permits the linear retained fraction in (II.77.6).

II.77.2 Optimizing the popularity threshold

The choice $\tau = \kappa/2$ makes the constants simple, but is not optimal for every downstream objective. To optimize the absolute bound in (II.77.4) using only $p \geq (\kappa - \tau)/(1 - \tau)$, maximize

$$F(\tau) = \frac{\tau^4(\kappa - \tau)}{1 - \tau}.$$

For $0 < \kappa < 1$, the maximizing root in $(0, \kappa)$ is

$$\tau_{\text{abs}} = \frac{5 + 3\kappa - \sqrt{25 - 34\kappa + 9\kappa^2}}{8}. \quad (\text{II.77.7})$$

Indeed, logarithmic differentiation gives $4/\tau - 1/(\kappa - \tau) + 1/(1 - \tau) = 0$, equivalently

$$4\tau^2 - (5 + 3\kappa)\tau + 4\kappa = 0.$$

The derivative changes sign from positive to negative at the smaller root. In particular $\tau_{\text{abs}}/\kappa \rightarrow 4/5$ as $\kappa \downarrow 0$. At $\kappa = 1$ the optimum is approached at the endpoint, and there is no need to use this formula in the proof of Corollary II.77.2.

More importantly, the threshold that optimizes a final Freiman restriction differs from the one that optimizes $|S - S|$ alone. Section II.80.4 gives its exact equation.

II.78 The sumset-growth input, with a complete proof

Source 24, delivered Section 7.

For self-containment, we record the classical Plünnecke–Ruzsa bound needed below. The short minimal-growth argument is due to Petridis [36]; no novelty is claimed for this section.

Lemma II.78.1 (Minimal growth). *Let A, B be nonempty finite subsets of an abelian group. Choose a nonempty $X \subseteq A$ minimizing $|X + B|/|X|$, and denote that minimum by K . For every finite T ,*

$$|X + B + T| \leq K|X + T|. \quad (\text{II.78.1})$$

Proof. The assertion is immediate for $T = \emptyset$. Otherwise order its elements $t_1, \dots, t_r$, and define

$$X_i = \{x \in X : x + t_i \notin \bigcup_{j < i} (X + t_j)\}.$$

The translates $X_i + t_i$ partition $X + T$, so $|X + T| = \sum_i |X_i|$. If $x \in X \setminus X_i$, then $x + t_i = x' + t_j$ for some $x' \in X$ and $j < i$. Adding any $b \in B$ shows that $((X \setminus X_i) + B) + t_i$ is already contained in $\bigcup_{j < i} (X + B + t_j)$. Thus the number of new points contributed by $X + B + t_i$ is at most

$$|X + B| - |(X \setminus X_i) + B| \leq K|X| - K|X \setminus X_i| = K|X_i|.$$

The inequality follows from minimality, with the empty-set case interpreted as $0 \geq K \cdot 0$. Summing over i proves (II.78.1). $\square$

Lemma II.78.2 (Ruzsa's triangle inequality). *For finite subsets U, V, W of an abelian group with $V \neq \emptyset$,*

$$|U - W||V| \leq |U - V||V - W|.$$

Proof. For each $d \in U - W$, choose a representation $d = u_d - w_d$. The map

$$(d, v) \mapsto (u_d - v, v - w_d)$$

from $(U - W) \times V$ to $(U - V) \times (V - W)$ is injective: the sum of the outputs recovers d , and then the first output recovers v . $\square$

Theorem II.78.3 (Plünnecke–Ruzsa in the required form). *If $|A + B| \leq C|A|$ for nonempty finite sets A, B , then for all nonnegative integers r, s ,*

$$|rB - sB| \leq C^{r+s}|A|. \quad (\text{II.78.2})$$

In particular, if $|\Gamma - \Gamma| \leq C|\Gamma|$, then

$$|r\Gamma - s\Gamma| \leq C^{r+s}|\Gamma|. \quad (\text{II.78.3})$$

Proof. Choose X as in Lemma II.78.1; then $K \leq C$. Applying that lemma with $T = (j - 1)B$ and inducting gives $|X + jB| \leq K^j|X|$ for every $j \geq 0$. Take $U = rB$, $V = -X$, and $W = sB$ in Lemma II.78.2. Then

$$|rB - sB||X| \leq |X + rB||X + sB| \leq K^{r+s}|X|^2.$$

Divide by $|X|$, and use $|X| \leq |A|$ and $K \leq C$. For (II.78.3), apply (II.78.2) with $A = \Gamma$ and $B = -\Gamma$, and interchange r, s if necessary. $\square$

All exponents in (II.78.3) count the *total* number of summands, positive and negative. In particular, $|2k\Gamma - 2k\Gamma|$ is controlled by $C^{4k}|\Gamma|$, not by $C^{2k}|\Gamma|$. The fibre argument below is designed to use a smaller sumset, rather than to miscount this exponent.

II.79 Vertical-fibre compression and Freiman restrictions

Source 24, delivered Section 8.

Let G, H be abelian groups, $B \subset G$ finite and nonempty, and $\phi : B \rightarrow H$. Its graph is

$$\Gamma = \{(x, \phi(x)) : x \in B\} \subset G \times H.$$

For an integer $k \geq 1$, define the vertical defect fibre

$$F_k = \{y \in H : (0, y) \in k\Gamma - k\Gamma\}, \quad Q_k = |F_k|. \quad (\text{II.79.1})$$

Then $0 \in F_k$, $F_k = -F_k$, and $Q_k \geq 1$. The condition $F_k = \{0\}$ is equivalent to ϕ already being a Freiman homomorphism of order k .

II.79.1 Counting only the relevant obstruction

Lemma II.79.1 (Compressed vertical-fibre estimate). *For every $k \geq 1$,*

$$Q_k |\Gamma| \leq |(k+1)\Gamma - k\Gamma|. \quad (\text{II.79.2})$$

Consequently, if $|\Gamma - \Gamma| \leq C|\Gamma|$, then

$$Q_k \leq C^{2k+1}. \quad (\text{II.79.3})$$

Proof. Use the map

$$F_k \times \Gamma \longrightarrow (k+1)\Gamma - k\Gamma, \quad (y, (x, \phi(x))) \longmapsto (x, \phi(x) + y).$$

Membership in the target follows from $(0, y) \in k\Gamma - k\Gamma$. The first coordinate of the image determines x , hence determines the graph point; its second coordinate then determines y . Thus the map is injective. Equation (II.79.3) follows from Theorem II.78.3 with $r = k+1$ and $s = k$. $\square$

The inspected repository proof of Lemma 7.5 uses a vertical fibre and bounds it through $2k\Gamma - 2k\Gamma$, resulting in C^{4k} ; its edited paper source explains why this repairs a larger-sumset problem in the printed proof [113]. Lemma II.79.1 uses the graph property more directly. One added graph point already makes the fibre injection work, so the total growth exponent becomes $2k+1$.

II.79.2 A prime-cyclic avoidance lemma with the exact symmetry gain

Lemma II.79.2 (Avoiding a symmetric defect set). *Let p be prime, let $F \subseteq \mathbb{F}_p$ be symmetric with $0 \in F$, and let $Q = |F| > 1$. For every integer $k \geq 1$, there are a nonzero $d \in \mathbb{F}_p$ and an integer $L \geq 0$ such that*

$$jd \notin F \quad (0 < |j| \leq kL), \quad \frac{L+1}{p} \geq \frac{1}{kQ}. \quad (\text{II.79.4})$$

Proof. Set

$$L = \left\lfloor \frac{p-2}{k(Q-1)} \right\rfloor. \quad (\text{II.79.5})$$

For each integer $1 \leq j \leq kL$ and $w \in F \setminus \{0\}$, the equation $jd = w$ has exactly one solution $d \in \mathbb{F}_p^\times$. Indeed, $j \leq kL \leq (p-2)/(Q-1) < p$, so j is invertible modulo p . There are at most $kL(Q-1) \leq p-2$ forbidden nonzero scalars, whereas $\mathbb{F}_p^\times$ has $p-1$ elements. Choose an allowed d . Positive multiples in this range avoid $F \setminus \{0\}$ and cannot be zero; negative multiples also avoid F because $F = -F$. This proves the first assertion, including the case $L = 0$.

For the second, integer division gives

$$k(Q-1)(L+1) \geq p-1.$$

Since $Q \leq p$,

$$\frac{L+1}{p} \geq \frac{p-1}{pk(Q-1)} \geq \frac{1}{kQ}.$$

The last comparison is exactly $Q(p-1) \geq p(Q-1)$. $\square$

The use of symmetry avoids counting positive and negative multiples as unrelated obstructions. Removing the zero element from the defect set also matters at small primes. These elementary details improve the constant in the extraction lemma, not just its power of C .

Theorem II.79.3 (Prime-codomain Freiman extraction). *Let B be a nonempty finite subset of an abelian group and $\phi : B \rightarrow \mathbb{F}_p$, with p prime. For every $k \geq 1$, there is $B' \subseteq B$ such that $\phi|_{B'}$ is a Freiman homomorphism of order k and*

$$|B'| \geq \frac{|B|}{kQ_k}. \quad (\text{II.79.6})$$

If the graph has $|\Gamma - \Gamma| \leq C|\Gamma|$, then

$$|B'| \geq \frac{|B|}{kC^{2k+1}}. \quad (\text{II.79.7})$$

When $Q_k = 1$, one can take $B' = B$.

Proof. The case $Q_k = 1$ follows from the definition of the defect fibre. Otherwise apply Lemma II.79.2 to F_k , obtaining d, L . Assign to each pair (x, j) with $x \in B$, $0 \leq j \leq L$, the colour

$$\phi(x) - jd \in \mathbb{F}_p.$$

Some colour a occurs at least $(L+1)|B|/p$ times. Since $L < p$ and $d \neq 0$, the same x cannot contribute twice to one colour. Hence

$$B' = \{x \in B : \phi(x) = a + j(x)d \text{ for some } 0 \leq j(x) \leq L\}$$

has at least $(L+1)|B|/p \geq |B|/(kQ_k)$ elements; the index $j(x)$ is unique.

Suppose $x_1 + \dots + x_k = x'_1 + \dots + x'_k$ with all entries in B' . The difference of the corresponding graph sums has the form

$$(0, Jd), \quad J = \sum_{i=1}^k j(x_i) - \sum_{i=1}^k j(x'_i), \quad |J| \leq kL.$$

Thus $Jd \in F_k$. By (II.79.4), J must be zero, which says exactly that the two output sums agree. Finally use $Q_k \leq C^{2k+1}$ from Lemma II.79.1. $\square$

[write] **6 October 2026; sources 21–27.** *Second route.* This is source 10's Theorem II.21.1 (the same set, $F_k = H_k$, the same injection into $(k+1)\Gamma - k\Gamma$ and the same avoidance length), whose prefactor $|B|/(k(h_k - 1) + 1)$ is at least $|B|/(kQ_k)$ (here $Q_k = h_k$); proved independently. The extension beyond prime cyclic domains is also in source 10's statement, which allows any abelian group.

Remark II.79.4 (Why the domain need not be prime cyclic). Primality was used only to invert j in the *codomain*. The domain was used only as an abelian group supporting equal sums and an injective graph projection. The theorem therefore genuinely extends the prime-cyclic-domain formulation in the source. For $k = 1$ every function already has the required property, so the quantitative extraction bound is deliberately nonoptimal in that trivial order.

II.79.3 Vector-space codomains in finite characteristic

Theorem II.79.5 (Linear hashing of the defect fibre). *Let $\phi : B \rightarrow \mathbb{F}_p^d$, where p is prime, and let Q_k be defined by (II.79.1). Put*

$$r = \left\lceil \log_p Q_k \right\rceil.$$

There is a subset $B' \subseteq B$ on which ϕ is a Freiman homomorphism of order k and

$$|B'| \geq \frac{|B|}{p^r} \geq \frac{|B|}{pQ_k}. \quad (\text{II.79.8})$$

If $|\Gamma - \Gamma| \leq C|\Gamma|$, it follows that $|B'| \geq |B|/(pC^{2k+1})$. For $Q_k = 1$, take $r = 0$ and $B' = B$.

Proof. Assume $Q_k > 1$. Choose a linear map $T : \mathbb{F}_p^d \rightarrow \mathbb{F}_p^r$ uniformly at random. For each fixed nonzero $y \in \mathbb{F}_p^d$, the image $T(y)$ is uniform, so $\Pr(T(y) = 0) = p^{-r}$. A union bound gives

$$\Pr(\ker T \cap (F_k \setminus \{0\}) \neq \emptyset) \leq (Q_k - 1)p^{-r} < 1.$$

Choose T for which $\ker T \cap F_k = \{0\}$. A fibre of $T \circ \phi$ has at least $|B|/p^r$ elements; let B' be such a fibre. If two k -term domain sums from B' agree, their output difference y lies in F_k . Its image under T is the difference between two sums of k copies of the same fibre value, hence zero. Therefore $y = 0$, proving the Freiman property. Since $p^{\lceil \log_p Q_k \rceil} \leq pQ_k$, the numerical bounds follow. $\square$

[write] **6 October 2026; sources 21–27.** *New in this report, and a comparison proved at the write.* Against source 10's torsion-budget Theorem II.21.3 on $V = \mathbb{F}_p^d$, $n = p^d$: for every multiple j of p , $V[j] = V$ and $jV = \{0\}$, so the summand $|V[j]| |F \cap jV| - 1$ of its budget equals $n - 1 > n - 2$; hence its budget fails as soon as $kL \geq p$, so $L \leq \lfloor (p-1)/k \rfloor$ and its retention $(L+1)/n$ is at most $(\lfloor (p-1)/k \rfloor + 1)p^{-d}$. For $k \geq 2$ this is less than p^{1-d} , while the bound here is $p^{-\lceil \log_p Q_k \rceil} \geq p^{1-d}$ whenever $Q_k \leq p^{d-1}$. So for $k \geq 2$ and $Q_k \leq p^{d-1}$ the bound here is strictly larger (the intake confirmed this numerically for $p \leq 11$, $d \in \{2, 3\}$, $k \in \{2, 3, 8\}$). For $d = 1$ source 10's prime-codomain bound is far better, as the remark below says, and in general neither dominates. Not formalized; the general finite abelian codomain remains open (item II.8).

This theorem complements the prime-cyclic interval argument rather than replacing it: for a large prime p and a small order k , Theorem II.79.3 has a better prefactor. For a vector-space codomain, the hashing bound is independent of its dimension. For a completely general finite abelian codomain, no corresponding uniform-characteristic assertion is being assumed.

II.80 Compiling energy bounds into coherent restrictions

Source 24, delivered Section 9.

Lemma II.80.1 (Quantitative compiler). *Suppose a graph Γ_0 on a domain of size m has a subgraph Γ_1 on B_1 satisfying*

$$|B_1| \geq a\gamma^u m, \quad |\Gamma_1 - \Gamma_1| \leq b\gamma^{-v} |\Gamma_1|,$$

where $a, b > 0$. If the codomain is $\mathbb{F}_p$, then an order- k Freiman restriction exists with size at least

$$\frac{a}{kb^{2k+1}} \gamma^{u+v(2k+1)} m. \quad (\text{II.80.1})$$

Proof. Apply Theorem II.79.3 to Γ_1 , then substitute the two assumed bounds. A subset of a graph is still a graph, with cardinality equal to that of its domain. $\square$

II.80.1 Proof of the headline theorem

The graph

$$\Gamma_0 = \{(x, \phi(x)) : x \in B_0\} \subset G \times \mathbb{F}_p$$

has m points. Its additive energy is precisely the number of respected quadruples, so is at least γm^3 . Corollary II.77.2, applied in the abelian product group, provides a graph subset Γ_1 on a domain B_1 with

$$|B_1| \geq 2^{-3}\gamma m, \quad |\Gamma_1 - \Gamma_1| \leq 2^{17}\gamma^{-6}|\Gamma_1|.$$

Use Lemma II.80.1 with $a = 2^{-3}$, $u = 1$, $b = 2^{17}$, and $v = 6$:

$$\begin{aligned} |B| &\geq \frac{2^{-3}}{k(2^{17})^{2k+1}} \gamma^{1+6(2k+1)} m \\ &= \frac{2^{-(34k+20)}}{k} \gamma^{12k+7} m. \end{aligned}$$

This proves Theorem F.29.1. For $k = 8$, the power of γ is $1 + 6 \cdot 17 = 103$, and the power-of-two loss is $3 + 17 \cdot 17 + 3 = 295$.

II.80.2 Separating the two sources of improvement

The following counterfactual calculations are useful for checking that no improvement is counted twice. The source's energy reduction retains $2^{-20}\gamma^{12}$ of the graph and has relative doubling $2^{58}\gamma^{-36}$; its extraction loss at order eight is $64C^{32}$ [1, 113].

Energy reduction	Freiman extraction	Order-eight retained fraction
Source	Source	$2^{-1882}\gamma^{1164}$
Source	Theorem II.79.3	$2^{-1009}\gamma^{624}$
Corollary II.77.2	Source	$2^{-553}\gamma^{193}$
Corollary II.77.2	Theorem II.79.3	$2^{-295}\gamma^{103}$

Table II.80.1: The improvement in the energy exponent is genuinely compositional: both the energy step and the fibre step contribute.

For example, using only the new extraction lemma gives $12 + 36 \cdot 17 = 624$ and $20 + 58 \cdot 17 + 3 = 1009$. Using only the new energy reduction gives $1 + 6 \cdot 32 = 193$ and $3 + 17 \cdot 32 + 6 = 553$. These equalities also make the provenance of the source's $1164 = 12 + 36 \cdot 32$ transparent.

II.80.3 An explicit corollary from a modern imported theorem

Theorem 1.2 of Reiher–Schoen [12] states that for $0 < \varepsilon < 1/2$, energy at least γm^3 yields a subset of size at least $(1 - \varepsilon)\gamma^{1/2}m$ with relative difference-set bound $2^{33}\varepsilon^{-9}\gamma^{-4}$. This theorem is an *external input* in the next corollary; its proof is not reproduced here.

Corollary II.80.2 (Modern-input restriction bound). *Under the hypotheses of Theorem F.29.1, for every $0 < \varepsilon < 1/2$ there is an order- k Freiman restriction of size at least*

$$\frac{(1 - \varepsilon)\varepsilon^{9(2k+1)}}{k 2^{33(2k+1)}} \gamma^{8k+9/2} m. \quad (\text{II.80.2})$$

In particular, taking $k = 8$ and $\varepsilon = 1/4$ gives

$$|B| \geq 3 \cdot 2^{-872} \gamma^{137/2} m. \quad (\text{II.80.3})$$

Proof. Apply the imported theorem to Γ_0 and use Lemma II.80.1 with $a = 1 - \varepsilon$, $u = 1/2$, $b = 2^{33}\varepsilon^{-9}$, and $v = 4$. For the displayed specialization, $b = 2^{51}$, $2k + 1 = 17$, and $(1 - \varepsilon)/k = 3/32$, giving $3 \cdot 2^{-5-51 \cdot 17} = 3 \cdot 2^{-872}$. $\square$

[write] **6 October 2026; sources 21–27.** *Second route.* This bound is identical to source 10's modern-input alternative (the paragraph after Corollary II.21.2), from the same imported Reiher-Schoen theorem and the same fibre argument.

The smaller energy exponent in (II.80.3) comes with a larger explicit constant loss. Consequently it does not uniformly dominate (F.29.2) for every $\gamma \in (0, 1]$. For fixed parameters one may take the larger of the two guaranteed sizes. This comparison illustrates why optimizing only an exponent, without its constants or the intended density regime, can mislead.

II.80.4 Optimizing for the final restriction rather than a single step

Use the unrounded popular-mass estimates in Theorem II.77.1. The resulting order- k retained fraction is at least

$$\frac{p}{2\sqrt{2}k} \left(2^{11} p^{-2} \tau^{-4} \right)^{-(2k+1)}. \quad (\text{II.80.4})$$

Its parameter-dependent part is $p^{4k+3} \tau^{8k+4}$. If the only information about p is its energy lower bound, set

$$a = 4k + 3, \quad b = 8k + 4, \quad H(\tau) = \left(\frac{\kappa - \tau}{1 - \tau} \right)^a \tau^b.$$

For $0 < \kappa < 1$, the unique maximizer in $(0, \kappa)$ is the smaller root of

$$b\tau^2 - [b(1 + \kappa) + a(1 - \kappa)]\tau + b\kappa = 0. \quad (\text{II.80.5})$$

Indeed, the logarithmic derivative vanishes exactly when

$$\frac{b}{\tau} - \frac{a}{\kappa - \tau} + \frac{a}{1 - \tau} = 0,$$

and the quadratic changes from positive to negative across its first root in the interval. As $\kappa \downarrow 0$,

$$\frac{\tau}{\kappa} \longrightarrow \frac{b}{a + b} = \frac{8k + 4}{12k + 7}. \quad (\text{II.80.6})$$

For order eight the limiting ratio is $68/103$, not $1/2$ and not the $4/5$ that optimizes the absolute difference bound. The clean dyadic theorem uses $\tau = \kappa/2$ to keep the statement simple; (II.80.4) and (II.80.5) supply an explicitly sharper parameter-specific certificate.

II.81 Sources 22, 25 and 27 on Lemma 15.4, and source 23 on Lemma 16.10 [write]

Added in the fifth write, 6 October 2026.

The next twenty-two sections continue two chapters printed earlier in this Part. *Sources 22, 25 and 27* (Sections II.82–II.102) return to Lemma 15.4, the count of degenerate arrangements that Lemma 15.5 must absorb, where source 18’s Section II.66 proved the sharp leading coefficient $k + C_d$. *Source 23’s Section II.103* returns to Lemma 16.10, where source 19’s chapter (Sections II.67–II.70) built common sampling and shared-slice interpolation. They are printed here, at the end of Part II, because no number printed earlier may move. Their introductions are Sections F.27, F.30, F.32 and F.28; their other sections are in Parts I, III and V. Their notation is in the table at the head of this Part.

First print, and independent re-derivations. The leading coefficient $k + C_d$ for $d \geq 2$ and primes $p > 2d$, and the sharp probability $p^{-2} + (p-1)p^{-2d}$ for one certificate, are source 18’s Theorem II.66.8 and Proposition II.66.3. Source 22 cites source 18 and reproves them; sources 25 and 27, written without access to source 18’s text (its archive sat unopened in `docs/incoming` at their pins, and source 27 has no pin), present them as their own: source 25’s Theorems F.30.1 and II.88.1 and source 27’s Theorems F.32.1, II.96.3 and Proposition II.94.3 are independent second and third routes. The same number appears as source 18’s C_d , source 22’s C_d , source 25’s H_d and source 27’s L_d ($C_8 = 2,585,444$).

What is new in this report. Source 22: a degree-stratified finite bound whose degree-one class count $J_{k,d}$ is polynomial in k ($30k^4 + 204k^3 + 111k^2 + 114k$ for $d = 2$) in place of source 18’s $J_{k,d}$ (Theorems II.83.2, II.84.1), and the case $d = 1$, which source 18 set aside: $\delta_{k,1}(p) = (k^2 + 3k)/p + O_k(p^{-2})$, exactly $4/p - 3/p^2$ for $k = 1$ (Theorem II.85.1). At Lemma 15.5 its thresholds gain nothing over source 18’s Corollary II.66.9 at leading order (note after Corollary II.84.5). Source 25: prime-power fields, the classification of the codimension-one components over an algebraically closed field, a genuine second-order lower bound (Theorem II.90.1), the small-characteristic counterexample (Example II.89.2), and a selection theorem for every $d \geq 2$ with exact survival probabilities (Theorem II.92.1). Source 27: the first order in every characteristic (Theorem II.96.3), the classification of the $S_{k,d}$ core directions with a cubic remainder (Theorem II.97.8), *the second coefficient for $d = 2$* , $(k^2 + 21k - 6)/2$, which answers source 18’s first question for $d = 2$ (Theorem F.32.3; general d reduced to one hyperplane statistic $b_{2,d,p}$, Corollary II.99.2), and the exact count for $k = 1$, $d = 2$ in characteristic at least 11 from a finite rational certificate (Section II.100). Its endpoint $k = d = 1$ (Proposition II.102.1) is source 22’s theorem at $k = 1$. Source 23: a localization to width at least $m/3$ (sharp for its cap), an integer lift of short progressions and an exact transport that partitions the full product without further width loss, a common base partition for many slice covers with width $L_m^A \geq (m/3)^A$, and affine gluing with $\sum_i Q_i + \sum_{i < j} Q_i Q_j$ candidates (Theorems II.103.6, II.103.7); it needs no inherited-coordinate hypothesis (source 19’s Corollary II.70.8 does) and is conditional on the slice covers it assumes. Its sampling lemma is measured against source 19’s theorem in the note after Lemma II.103.8. Lemma 16.10 itself is not proved by any source; its closing comparison is still open.

Relation to the formal statements (checked at HEAD 2dfe46f44). `lemma_15_4` and `lemma_15_5` (Sections14_15.lean, lines 402, 414) have exact companions (`Proofs15DegenerateArrangements.lean`, line 1131; `Proofs15Restriction.lean`, line 3086). The degeneracy predicate `GeneralArrangement.IsDegenerate` (line 354) is the one of sources 22, 25 and 27 (all moments, the modular-multiple exclusion), and every Lean name they use exists: `GeneralArrangement` (line 192), `walshCharacter` (line 310), `IsTernaryCoefficient` (line 319), `IsModularMultiple` (line 324), `signedWalshMoment` (line 329), `arrangementParityCoefficient` (line 341), `arrangementMoment` (line 346), `degenerateGeneralArrangementCount` (line 364), `arrangementSelectionExponent` (line 409). `lemma_16_10` (Section16.lean, line 851) is open. Source 23’s widths compare with the formal `Box.short_parent_partition` (`Proofs16ShortParents.lean`, line 23: width $m/8$), `box_`

product_tiling_of_contained_axis (Proofs16RelativeSteps.lean, line 89: $v^2 \leq \text{width} - 1$, final lengths $v - 1$ or v) and section16_local_affine_lift_sqrt (Proofs16LiftWidth.lean, line 31: width $\sqrt{(m/8)^a}/4$); notes after its Lemma II.103.1 and Theorem II.103.6. Nothing in these sections is formalized.

II.82 Arrangements, parity, and the degree of a certificate

Source 22, delivered Section 2.

In Sections II.82–II.84, let $k \geq 1$, $d \geq 2$, $q = 2d$, and let $p > 4d$ be prime. Here q is the number of cubes, not a field cardinality. Later sections reuse q for the customary finite-field cardinality, after saying so explicitly. Set

$$s = (\underbrace{1, \dots, 1}_d, \underbrace{-1, \dots, -1}_d), \quad H = \{r \in \mathbb{F}_p^q : s \cdot r = 0\}.$$

An arrangement is parametrized by

$$h \in \mathbb{F}_p^k, \quad r \in H, \quad (y^1, \dots, y^q) \in (\mathbb{F}_p^k)^q,$$

and its labeled vertices are

$$\rho(\epsilon, j) = (y_1^j + \epsilon_1 h_1, \dots, y_k^j + \epsilon_k h_k, r_j), \quad \epsilon \in \{0, 1\}^k, \quad j \in [q]. \quad (\text{II.82.1})$$

There are p^D parameterized arrangements, where $D = (q + 1)k + q - 1$. All probabilities use the uniform measure on this parameter space.

Definition II.82.1. A certificate is a function $\eta : \{0, 1\}^k \times [q] \rightarrow \{-1, 0, 1\}$. It succeeds when

$$\sum_{\epsilon, j} \eta(\epsilon, j) \prod_{i \in A} \rho(\epsilon, j)_i = 0 \quad \text{for every } A \subseteq [k + 1]. \quad (\text{II.82.2})$$

Write $\pi(\epsilon) = (-1)^{|\epsilon|}$ and $\eta_0(\epsilon, j) = s_j \pi(\epsilon)$. The arrangement is degenerate if a successful certificate is not in $\text{span}\{\eta_0\}$. Its degeneracy probability is denoted $\delta_{k,d}(p)$.

The universal certificate η_0 always succeeds: parity annihilates proper horizontal monomials; the remaining two tests are proportional to $\sum s_j = 0$ and $s \cdot r = 0$. A zero side identifies two labeled vertices and gives another successful certificate. Thus the exact probability of this first source of degeneracy is

$$z_k(p) = 1 - (1 - p^{-1})^k. \quad (\text{II.82.3})$$

For the rest of this section fix nonzero sides.

II.82.1 Removing the parity factor

Normalize a certificate by

$$v_j(\epsilon) = \eta(\epsilon, j) / \eta_0(\epsilon, j) = s_j \pi(\epsilon) \eta(\epsilon, j).$$

Each v_j still takes values in $\{-1, 0, 1\}$. Let its unique multilinear interpolation polynomial over $\mathbb{F}_p$ be

$$v_j(\epsilon) = \sum_{T \subseteq [k]} a_{j,T} \epsilon^T, \quad \epsilon^T = \prod_{i \in T} \epsilon_i.$$

Define the degree of η to be $t = \max_j \deg v_j$, taking a nonzero constant slice to have degree zero. The zero certificate is excluded whenever an equivalence class is counted.

Lemma II.82.2 (Moment–degree duality). *Put $c_{j,S} = \sum_{\epsilon} \eta(\epsilon, j) \epsilon_S$. If a nonzero certificate has degree t , then the smallest $|S|$ for which some $c_{j,S}$ is nonzero is exactly $k - t$. In particular, degree zero is precisely the pure-parity case.*

Proof. Parity summation of a Boolean monomial is zero unless that monomial uses all k coordinates, when it is $(-1)^k$. Hence

$$c_{j,S} = s_j (-1)^k \sum_{T: T \cup S = [k]} a_{j,T}. \quad (\text{II.82.4})$$

If $|S| < k - t$, no set T of size at most t can occur in the sum. Conversely, choose $a_{j,T} \neq 0$ with $|T| = t$ and set $S = [k] \setminus T$. The only possibly nonzero summand is $a_{j,T}$, proving equality. $\square$

II.82.2 The rank gained at a given degree

The following lemma reproves and preserves the degree information in the paired-moment argument of [105].

Lemma II.82.3 (Degree-sensitive certificate probability). *Let a nonparity certificate have degree $t \geq 1$. Conditional on any fixed nonzero sides, its probability of success is at most*

$$Q_t(p) = \begin{cases} p^{-2} + (p-1)p^{-q}, & t = 1, \\ p^{-t-1}, & 2 \leq t \leq k. \end{cases} \quad (\text{II.82.5})$$

The degree-one bound is attained by an individual certificate.

Proof. Choose $B \subseteq [k]$ of size $k - t$ with $b = (c_{1,B}, \dots, c_{q,B}) \neq 0$. For $A \subseteq [k]$ define

$$P_{j,A}(y^j) = \sum_{\epsilon} \eta(\epsilon, j) \prod_{i \in A} (y_i^j + \epsilon_i h_i).$$

Expansion and the minimality of $|B|$ give

$$P_{j,A} = \sum_{S \subseteq A} h_S c_{j,S} y_{A \setminus S}^j, \quad P_{j,B} = h_B b_j.$$

The tests for B and $B \cup \{k+1\}$ force $\sum_j b_j = 0$ and $b \cdot r = 0$.

For each $i \notin B$, the equation $\sum_j P_{j,B \cup \{i\}} = 0$ is affine in the horizontal variables. Choose j_0 with $b_{j_0} \neq 0$ and hold all variables other than $y_i^{j_0}$, $i \notin B$, fixed. Their coefficient matrix is diagonal with nonzero diagonal $h_B b_{j_0}$. Thus these equations have rank t . If b is not proportional to s , then $b \cdot r = 0$ is an independent restriction on H , giving probability at most $p^{-1}p^{-t}$. An inconsistent condition $\sum b_j \neq 0$ only reduces the probability.

If $b = cs$ with $c \neq 0$, pair each affine equation with its vertical counterpart $\sum_j r_j P_{j,B \cup \{i\}} = 0$. For $r_a \neq r_b$, the columns indexed by y_i^a, y_i^b have coefficient block

$$h_B c \begin{pmatrix} s_a & s_b \\ r_a s_a & r_b s_b \end{pmatrix},$$

with nonzero determinant. Distinct $i \notin B$ give disjoint diagonal blocks, so the rank is at least $2t$ when r is not constant. It remains at least t when r is constant. There are p constant vectors among the p^{q-1} vectors of H . The probability is therefore at most

$$(1 - p^{-(q-2)})p^{-2t} + p^{-(q-2+t)}. \quad (\text{II.82.6})$$

For $t = 1$ this equals the first expression in (II.82.5). For $t \geq 2$, its ratio to p^{-t-1} is at most $p^{1-t} + p^{3-q} \leq 2/p < 1$, since $q \geq 4$.

Finally, fix an index i and take $\eta = s_j \pi(\epsilon) \mathbf{1}_{\{\epsilon_i=0\}}$. Only the equations $\sum_j s_j y_i^j = 0$ and $\sum_j s_j r_j y_i^j = 0$ remain nonautomatic. Their ranks are exactly two or one according as r is nonconstant or constant. This realizes the first line of (II.82.5). $\square$

II.82.3 Equivalent certificates and cumulative counts

Two certificates related by $\eta' = a\eta + b\eta_0$, $a \neq 0$, have the same success locus. In normalized form this is the affine change $v' = av + b\mathbf{1}$. The count of these equivalence classes is elementary.

Lemma II.82.4 (Ternary affine classes). *For $p \geq 5$, nonconstant ternary patterns on n positions have*

$$A_n = \frac{3^n - 2^{n+1} + 1}{2} \quad (\text{II.82.7})$$

classes under invertible affine changes of values.

Proof. Two-valued patterns correspond to unordered nontrivial bipartitions, giving $(2^n - 2)/2$ classes. A pattern using all three values can only be affinely equivalent to itself or its negative: an affine permutation of $\{-1, 0, 1\}$ has translation part zero, by summing the three values, and multiplier 1 or -1 . Inclusion-exclusion counts $3^n - 3 \cdot 2^n + 3$ three-valued patterns. Half this number, added to the two-valued count, gives (II.82.7). $\square$

Let C_t be the actual number of nonconstant normalized certificate classes of degree at most t ; they need not succeed. Set

$$H_t(k) = \sum_{j=0}^t \binom{k}{j}.$$

A degree-at-most- t multilinear polynomial is determined by its values on the Boolean vectors of weight at most t . Indeed its coefficients are recovered successively by Boolean interpolation. Restriction to these vectors in every slice therefore injects affine classes. Consequently

$$C_t \leq A_q H_t(k), \quad C_0 = A_q, \quad C_k = A_{q2^k}. \quad (\text{II.82.8})$$

This inequality is valid over $\mathbb{F}_p$ directly; it does not require identifying a finite-field degree with a real polynomial degree.

II.83 An exact polynomial count for the first nonparity degree

Source 22, delivered Section 3.

II.83.1 Affine maps with three allowed values

Lemma II.83.1 (Affine ternary slices). *Over $\mathbb{F}_p$, $p \geq 5$, an affine function on $\{0, 1\}^k$ with values in $\{-1, 0, 1\}$ is one of the following:*

1. a constant;

2. a function with one slope ± 1 and one of two intercepts;
3. a function with one slope ± 2 and its unique intercept;
4. a function with two slopes, both in $\{-1, 1\}$, and their unique admissible intercept.

Hence the numbers of ternary and binary affine functions are respectively

$$f_k = 2k^2 + 4k + 3, \quad b_k = 2k + 2. \quad (\text{II.83.1})$$

Proof. Fix an active slope a and put $S = \{-1, 0, 1\}$. The values on the remaining coordinates must lie in $S \cap (S - a)$. This intersection is empty unless $a \in \{\pm 1, \pm 2\}$; it has two elements for $a = \pm 1$ and one for $a = \pm 2$. An affine Boolean map into a two-point set has at most one active coordinate: two active slopes would force the four values $x, x+b, x+c, x+b+c$ into that set, implying $b = c$ and $2b = 0$. This is impossible in odd characteristic. Thus at most two coordinates are active. The listed intercepts follow by checking their two or four vertex values. Counting gives $3 + 6k + 4\binom{k}{2} = f_k$. A binary affine map is constant or a coordinate or its complement, giving b_k . $\square$

Applying the orbit argument of Lemma II.82.4 to all q slices gives the exact cumulative count

$$C_1 = U_1 = \frac{f_k^q - 2b_k^q + 1}{2}. \quad (\text{II.83.2})$$

This already replaces a count exponential in 2^k by a polynomial in k of degree $2q$. The success equations impose further useful restrictions.

II.83.2 Balancing the slopes

Write a degree-one slice as $v_j = c_j + \sum_i a_{ij}\epsilon_i$. Equation (II.82.4), applied with $S = [k] \setminus \{i\}$, shows that success requires

$$\sum_{j=1}^q s_j a_{ij} = 0 \quad (1 \leq i \leq k). \quad (\text{II.83.3})$$

The slopes have their integer representatives in $\{-2, -1, 0, 1, 2\}$. Since $p > 2q = 4d$, their sums in (II.83.3) vanish over the field exactly when the corresponding integer sums vanish.

Introduce independent Laurent variables $z_1, \dots, z_k$ and set $u_i = z_i + z_i^{-1}$. Define

$$F_k(z) = 3 + \sum_{i=1}^k (2u_i + z_i^2 + z_i^{-2}) + \sum_{i < j} u_i u_j, \quad (\text{II.83.4})$$

$$B_k(z) = 2 + \sum_{i=1}^k u_i. \quad (\text{II.83.5})$$

The coefficient of z^a in F_k counts the allowed intercepts for slope vector a . The corresponding statement for binary slices uses B_k . Both Laurent polynomials are invariant under inversion of all variables.

Theorem II.83.2 (Balanced degree-one certificate count). *For $p > 4d$, the exact number of degree-one normalized certificate classes satisfying all necessary slope balances is*

$$J_{k,d} = \frac{\text{CT } F_k^{2d} - 2 \text{CT } B_k^{2d} - 3^{2d} + 2^{2d+1}}{2}. \quad (\text{II.83.6})$$

The constant term is taken in all k variables. For fixed d , $J_{k,d}$ is a polynomial in k of degree exactly $2d$, with leading coefficient

$$\lambda_d = 2^{2d-1}(2d)![z^{2d}] \frac{e^{-z/2}}{\sqrt{1-z}} > 0. \quad (\text{II.83.7})$$

In particular,

$$J_{k,2} = 30k^4 + 204k^3 + 111k^2 + 114k. \quad (\text{II.83.8})$$

This counts potentially successful classes, not distinct geometric success loci and not actual successful certificates in one arrangement.

Proof. The first d and last d slices contribute opposite slope sums. Their balanced tuples are counted by $\text{CT } F_k(z)^d F_k(z^{-1})^d = \text{CT } F_k(z)^{2d}$. Similarly, the binary balanced tuples number $\text{CT } B_k^{2d}$. Slope balance is preserved by affine value changes, so the two-value and three-value orbit argument gives

$$\frac{\text{CT } F_k^{2d} - 2 \text{CT } B_k^{2d} + 1}{2}$$

classes of degree at most one satisfying balance. Every pure-parity class satisfies balance, and their number is A_{2d} . Subtracting it proves (II.83.6).

For polynomiality put $q = 2d$. A constant-term monomial in F_k^q contains each active variable in at least two slice factors: one nonzero signed exponent cannot sum to zero. Each slice contains at most two active variables. Thus at most q variables are active. Classify terms by their actual active set. For every $0 \leq j \leq q$ the contribution is a constant, independent of k , times $\binom{k}{j}$. Therefore $\text{CT } F_k^q$ is a polynomial of degree at most q in k . The same argument gives degree at most $\lfloor q/2 \rfloor$ for $\text{CT } B_k^q$.

In the top-degree contribution, each of the q slices uses exactly two different variables, with slopes ± 1 , and each of the q active variables occurs twice. Represent the slices by labeled vertices and each active variable by an edge between its two slices. The result is a loopless 2-regular multigraph; a two-vertex component has a pair of parallel edges. Each edge has two balanced sign choices. After assigning distinct variables to edges, a multigraph with edge multiplicities m_{ab} has weight $1/\prod m_{ab}!$. The weighted exponential generating function for these graphs is

$$\exp \left(\sum_{j \geq 2} \frac{z^j}{2j} \right) = \frac{e^{-z/2}}{\sqrt{1-z}}.$$

For $j \geq 3$ this uses the $(j-1)!/2$ undirected cycles; for $j = 2$ the parallel pair has weight $1/2$, giving the same term $z^2/4$. The leading coefficient of $\text{CT } F_k^q$ is consequently $2^q q! [z^q] e^{-z/2} (1-z)^{-1/2}$. The binary correction has lower degree. Dividing by two proves (II.83.7) and positivity.

For $q = 4$, expansion gives

$$\begin{aligned} \text{CT } F_k^4 &= 60k^4 + 408k^3 + 246k^2 + 312k + 81, \\ \text{CT } B_k^4 &= 12k^2 + 42k + 16. \end{aligned}$$

Substitution in (II.83.6) gives (II.83.8). These expansions can be obtained by classifying active sets of size at most four, or by the exact moment calculation below. $\square$

II.83.3 An efficient exact evaluation of the constant terms

The displayed constant-term formula does not require expanding in k Laurent variables. Set $S = \sum_i u_i$ and $T = \sum_i u_i^2$. Then

$$F_k = 3 - 2k + 2S + \frac{S^2 + T}{2}, \quad B_k = 2 + S,$$

and

$$\text{CT } u_i^a = \begin{cases} \binom{a}{a/2}, & a \text{ even,} \\ 0, & a \text{ odd.} \end{cases}$$

Independence of the Laurent variables yields the recurrence

$$M_{k+1}(a, b) = \sum_{i=0}^a \sum_{j=0}^b \binom{a}{i} \binom{b}{j} M_k(a-i, b-j) \mu_{i+2j}, \quad (\text{II.83.9})$$

where $M_k(a, b) = \text{CT } S^a T^b$, $\mu_r = \text{CT}(z + z^{-1})^r$, and $M_0(a, b) = \mathbf{1}_{\{a=b=0\}}$. Only $a + 2b \leq 2q$ is needed. The script `code/verify_certificates.py` implements this recurrence using integers and compares it with independent sparse Laurent convolution.

II.84 The finite exceptional-set bound and its quantitative use

Source 22, delivered Section 4.

II.84.1 The degree-zero hyperplanes

For completeness, the earlier coefficient C_d has the following exact interpretation. A pure-parity certificate has $\eta(\epsilon, j) = c_j \pi(\epsilon)$ with $c_j \in \{-1, 0, 1\}$. For nonzero sides it succeeds exactly when

$$\sum_j c_j = 0, \quad c \cdot r = 0.$$

Excluding the universal direction $\text{span}\{s\}$, these equations give C_d distinct hyperplanes on H .

Here is the count. There are T_{2d} integer-balanced ternary vectors c . Write $v_j = s_j c_j$. Two vectors define the same hyperplane on H exactly when $v' = av + b\mathbf{1}$. The admissible two-valued patterns correspond to balanced subsets of the d positive and d negative positions, of which there are $\binom{2d}{d}$ including the empty and full subsets. They contribute $(\binom{2d}{d} - 2)/2$ classes. The admissible three-valued patterns number $T_{2d} - 3\binom{2d}{d} + 3$ and are paired by negation. Adding gives (F.27.1). The assumption $p > 2d$ prevents integer balance from aliasing modulo p and $p \geq 5$ validates the orbit count.

For reference $C_2 = 4$, $C_3 = 51$, $C_4 = 484$, and $C_8 = 2\,585\,444$. Their union has probability at most C_d/p . Any two distinct such hyperplanes intersect in codimension two in H .

II.84.2 Summing the degree strata correctly

Define an explicit cumulative profile by

$$U_0 = A_q, \quad U_1 = \frac{f_k^q - 2b_k^q + 1}{2}, \quad U_t = A_{qH_t(k)} \quad (2 \leq t \leq k). \quad (\text{II.84.1})$$

For $k = 1$ only the first two terms are used. This profile is increasing, $C_t \leq U_t$, and equality holds for $t = 0, 1, k$.

Theorem II.84.1 (Stratified finite degeneracy bound). *Let $k \geq 1$, $d \geq 2$, and $p > 4d$ be prime. Then*

$$\delta_{k,d}(p) \leq z_k(p) + (1 - z_k(p)) \left[\frac{C_d}{p} + J_{k,d}(p^{-2} + (p-1)p^{-2d}) + \sum_{t=2}^k \frac{U_t - U_{t-1}}{p^{t+1}} \right]. \quad (\text{II.84.2})$$

The empty sum is zero. One may take the minimum of this expression and one. Multiplication by p^D gives the corresponding count of arrangements.

Proof. Condition on nonzero sides. The pure-parity classes contribute at most C_d/p . Among degree-one classes, only the $J_{k,d}$ balanced ones can succeed, and Lemma II.82.3 applies to each. For degree $t \geq 2$, that lemma gives an upper bound

$$\sum_{t=2}^k (C_t - C_{t-1})w_t, \quad w_t = p^{-t-1}.$$

For $k \geq 2$, summation by parts rewrites this as

$$C_k w_k - C_1 w_2 + \sum_{t=2}^{k-1} C_t (w_t - w_{t+1}).$$

The differences $w_t - w_{t+1}$ are positive, and $C_1 = U_1$. Replacing the other C_t by their cumulative upper bounds and reversing the summation by parts gives $\sum_{t=2}^k (U_t - U_{t-1})w_t$. The case $k = 1$ has no such tail. Finally include the zero-side event (II.82.3). $\square$

[write] **6 October 2026; sources 21–27.** *Against source 18 and the formal statement* (checked at the write). With $q = 2d$, the probability $Q_1(p) = p^{-2} + (p-1)p^{-2d}$ of Lemma II.82.3 is source 18's $Q_d(p)$, and the remark at the end of this section shows that the right side of (II.84.2) is at most the bound obtained by charging $Q_1(p)$ to every nonparity class, which is source 18's finite bound (II.66.11). That bound lies below the formal $k3^{2d2^k}/p$ of lemma_15_4 for $d \geq 2$ and $p > 2d$ (note after Theorem II.66.8). So this theorem implies the formal bound for $d \geq 2$ and primes $p > 4d$; it is not formalized. It needs $p > 4d$ where source 18 needs $p > 2d$.

Remark II.84.2 (Why cumulative differences need this proof). It is not asserted that $C_t - C_{t-1} \leq U_t - U_{t-1}$ for every t . Such an assertion would not follow from cumulative bounds. The decreasing weights and exact endpoint at degree one are essential to the valid aggregate comparison. Also, $J_{k,d}$ replaces only the degree-one term; its use does not silently subtract unbalanced classes from the higher-degree profile.

Corollary II.84.3 (A smaller second-order upper coefficient). *For fixed $k, d \geq 1$ with $d \geq 2$, as $p \rightarrow \infty$ through primes,*

$$\delta_{k,d}(p) \leq \frac{k + C_d}{p} + \frac{J_{k,d} - kC_d - \binom{k}{2}}{p^2} + O_{k,d}(p^{-3}). \quad (\text{II.84.3})$$

In particular the positive contribution to this upper coefficient from nonparity certificates is polynomial of degree $2d$ in k .

Proof. Expand $z_k = k/p - \binom{k}{2}/p^2 + O_k(p^{-3})$. The extra term $(p-1)p^{-2d}$ is $O_d(p^{-3})$ and every tail term has order at most p^{-3} . Multiplication in (II.84.2) gives the result. $\square$

[write] **6 October 2026; sources 21–27.** *An upper coefficient, now known exactly for $d = 2$.* Source 27’s Theorem F.32.3 determines the second coefficient for $d = 2$: $(k^2 + 21k - 6)/2$, which is 8 at $k = 1$ and 20 at $k = 2$, against the upper coefficients $J_{1,2} - 4 = 455$ and $J_{2,2} - 8 - 1 = 2,775$ given here. As the text says, the class count alone cannot give the true coefficient.

The true second coefficient need not equal the coefficient in (II.84.3): hyperplane intersections and coincident nonparity loci still matter. Retaining the zero-side and parity hyperplanes gives the earlier lower estimate

$$\delta_{k,d}(p) \geq \frac{k + C_d}{p} - \frac{\binom{k+C_d}{2}}{p^2}.$$

Together these bounds recover the sharp leading term while narrowing the source of its next-order error.

II.84.3 Examples and the size of the improvement

Let $J_{k,d}^{\text{old}} = A_{2d2^k} - A_{2d}$ be the earlier count of all nonparity classes. The following table compares coefficients that had previously been charged at the same p^{-2} scale. The new finite theorem also includes its explicitly displayed higher-degree tail.

k	d	$J_{k,d}$	$\log_{10}(U_1 - U_0)$	$\log_{10} J_{k,d}^{\text{old}}$
1	2	459	3.477	3.477
2	2	2 784	4.805	7.332
4	2	22 968	6.528	30.235
8	2	235 344	8.548	488.271
2	8	$1.68440 \cdot 10^{18}$	20.159	30.235
4	8	$4.58798 \cdot 10^{23}$	27.020	121.842
8	8	$3.27723 \cdot 10^{29}$	35.094	1953.988

Rounded values in this table are explanatory only. Exact integers are recorded in `data/certificate_verification.json`.

For example, for $k = d = 2$, one has

$$C_d = 4, \quad J_{2,2} = 2784, \quad U_1 - U_0 = 63\,840, \quad U_2 - U_1 = 21\,393\,960.$$

Thus the conditional exceptional probability is bounded by

$$\frac{4}{p} + 2784 \left(\frac{1}{p^2} + \frac{p-1}{p^4} \right) + \frac{21\,393\,960}{p^3}. \quad (\text{II.84.4})$$

One sees explicitly why a finite-size comparison must retain the tail: the asymptotic gain can be much larger than the gain at moderate p .

II.84.4 Explicit absorption thresholds

For $0 < \varepsilon \leq 1$ define

$$\mathcal{P}_{k,d}(\varepsilon) = \max \left\{ 4d, \frac{(k+1)(k+C_d)}{\varepsilon}, \left(\frac{2(k+1)\mathbf{J}_{k,d}}{\varepsilon} \right)^{1/2}, \max_{2 \leq t \leq k} \left(\frac{(k+1)(\mathbf{U}_t - \mathbf{U}_{t-1})}{\varepsilon} \right)^{1/(t+1)} \right\}. \quad (\text{II.84.5})$$

Omit the inner maximum when $k = 1$.

Corollary II.84.4 (Finite exceptional-set absorption). *Every prime $p > \mathcal{P}_{k,d}(\varepsilon)$ satisfies $\delta_{k,d}(p) \leq \varepsilon$.*

Proof. Use $z_k \leq k/p$, $1 - z_k \leq 1$, and $p^{-2} + (p-1)p^{-2d} \leq 2p^{-2}$. The resulting sum has $k+1$ nonnegative terms: one leading term, one degree-one term, and $k-1$ higher-degree terms. Each is at most $\varepsilon/(k+1)$ under (II.84.5). $\square$

The equal allocation in (II.84.5) is convenient, but it need not be the smallest available sufficient condition. There is a second explicit form retaining the optimal leading scale as $\varepsilon \downarrow 0$. Write $a = k + C_d$, $J = \mathbf{J}_{k,d}$, and $\Delta_t = \mathbf{U}_t - \mathbf{U}_{t-1}$, and set

$$\mathcal{A}_{k,d}(\varepsilon) = \frac{a}{\varepsilon} + \frac{J}{a} + \frac{J\varepsilon^{2d-3}}{a^{2d-2}} + \sum_{t=2}^k \frac{\Delta_t \varepsilon^{t-1}}{a^t}. \quad (\text{II.84.6})$$

Corollary II.84.5 (A threshold with the sharp leading scale). *If $p > 4d$ is prime and $p \geq \mathcal{A}_{k,d}(\varepsilon)$, then $\delta_{k,d}(p) \leq \varepsilon$. For fixed k, d ,*

$$\mathcal{A}_{k,d}(\varepsilon) = \frac{k + C_d}{\varepsilon} + \frac{\mathbf{J}_{k,d}}{k + C_d} + O_{k,d}(\varepsilon).$$

Proof. Theorem II.84.1 gives

$$p\delta_{k,d}(p) \leq a + \frac{J}{p} + \frac{J}{p^{2d-2}} + \sum_{t=2}^k \frac{\Delta_t}{p^t}.$$

Since $p \geq a/\varepsilon$, replace $1/p$ on the right by ε/a . The resulting right side is exactly $\varepsilon\mathcal{A}_{k,d}(\varepsilon)$, at most εp . The asymptotic expansion follows from $d \geq 2$. $\square$

[write] **6 October 2026; sources 21–27.** At Lemma 15.5 there is no gain at leading order (checked at the write). There the scale is $\varepsilon = \lambda = (\alpha\eta/4)^{E_k} \beta^{15}$, astronomically small, so a threshold is judged by its leading term in $1/\lambda$. Source 18’s Corollary II.66.9 asks for $\lambda p \geq k + 2,585,445$, i.e. $p \geq (k + C_8 + 1)/\lambda$, or, by its quadratic form (II.66.16), $p \geq (k + C_8)/\lambda + (18/17)J_k/(k + C_8) + o(1)$. The threshold $\mathcal{A}_{k,8}(\lambda) = (k + C_8)/\lambda + \mathbf{J}_{k,8}/(k + C_8) + O(\lambda)$ of this corollary has the same leading term; its lower-order term is smaller, because $\mathbf{J}_{k,8}$ is polynomial in k while source 18’s J_k grows like $3^{2^{k+4}}$, but that term is negligible against $1/\lambda$. The equal-allocation threshold (II.84.5) is worse: $\mathcal{P}_{k,8}(\lambda) \geq (k+1)(k+C_8)/\lambda$, larger by the factor $k+1$. So the sentence below that Lemma 15.4 “can therefore be replaced, for that purpose, by the concrete condition $p > \mathcal{P}_{k,8}(\varepsilon)$ ” is correct, and the replacement is weaker than source 18’s at leading order; the text says only that the equal-budget threshold “is not claimed to dominate every earlier optimized threshold”. The finite bound of Theorem II.84.1 at moderate p is where this section improves on source 18.

Either sufficient condition may be used. More precisely, the complete right side of (II.84.2) is uniformly no larger than the earlier bound that assigns $Q_1(p)$ to every nonparity class: all the tail weights are at most $Q_1(p)$, their profile differences sum to $U_k - U_1$, and $J \leq U_1 - U_0$. The convenient equal-budget threshold (II.84.5) by itself is not claimed to dominate every earlier optimized threshold.

In Gowers' Lemma 15.5, and in the corrected transcription, the random selection step produces a scale

$$\varepsilon = (\alpha\eta/4)^{2^{2^{k+4}+k+3}}\beta^{15}, \quad 0 < \alpha, \eta, \beta \leq 1.$$

The use of Lemma 15.4 at that step is to ensure that at most εp^{17k+15} eight-arrangements are degenerate. It can therefore be replaced, for that purpose, by the concrete condition

$$p > \mathcal{P}_{k,8}(\varepsilon).$$

This changes only the exceptional-set absorption requirement. All other hypotheses of the selection lemma remain in force. In particular it does not establish the downstream partition and inverse-theorem estimates needed to state a new bound for Szemerédi's theorem. The substantial point is that high-degree certificate counts enter through roots of orders $t + 1$, instead of all being charged through a square root or a single worst-case term.

II.85 The exceptional two-cube case: a signed-coordinate arrangement

Source 22, delivered Section 5.

The assumption $d \geq 2$ in the codimension-two bound for nonparity certificates cannot be discarded: when $d = 1$, the vertical constraint is $r_1 = r_2$, so the vertical equations supply no independent information. Nevertheless the first-order count can be determined exactly.

Theorem II.85.1 (Sharp first-order degeneracy for two cubes). *Let $k \geq 1$ and let $p > \max\{3, 2k+1\}$ be prime. Define*

$$L_k(p) = 1 - (1 - p^{-1})^k \prod_{i=1}^k \left(1 - \frac{2i+1}{p}\right), \quad R_k = \frac{1}{2} \left(3^{2^k} - (2k^2 + 4k + 3)\right)^2.$$

For the proportion $\delta_{k,1}(p)$ of degenerate 1-arrangements,

$$L_k(p) \leq \delta_{k,1}(p) \leq L_k(p) + (1 - p^{-1})^k R_k p^{-2}.$$

Consequently

$$\delta_{k,1}(p) = \frac{k^2 + 3k}{p} + O_k(p^{-2}).$$

For $k = 1$ the error term in the finite inequality is zero, and hence $\delta_{1,1}(p) = 4p^{-1} - 3p^{-2}$ for every prime $p \geq 5$.

Proof. The proportion of arrangements having a zero side is exactly $1 - (1 - p^{-1})^k$, and every such arrangement is degenerate. Condition on nonzero sides, write $z_i^j = y_i^j/h_i$, and set $t_i = z_i^1 - z_i^2$. The

variables $t_1, \dots, t_k$ are independent and uniform in $\mathbb{F}_p$. For a slice coefficient pattern $\eta : \{0, 1\}^k \rightarrow \{-1, 0, 1\}$, put

$$F_\eta(X) = \sum_{\epsilon \in \{0, 1\}^k} \eta(\epsilon) \prod_{i=1}^k (X_i + \epsilon_i).$$

The map $\eta \mapsto F_\eta$ is an invertible linear transformation: in one coordinate its basis is $X, X + 1$, and the general transform is the tensor product of these invertible two-dimensional transforms. Comparing coefficients of $X_{[k] \setminus A}$ shows that all the horizontal moment equations are equivalent to

$$F_{\eta_1}(X + t) + F_{\eta_2}(X) = 0. \quad (\text{D1.1})$$

The vertical equations duplicate them because $r_1 = r_2$.

We first describe the solution set of (D1.1) for fixed patterns. For any multiaffine polynomial F , let

$$D_u = \sum_i u_i \frac{\partial}{\partial X_i}.$$

Since $\deg F \leq k < p$, Taylor's identity gives $F(X + u) = \exp(D_u)F(X)$. On the space of polynomials of degree at most k , the operator D_u is nilpotent. Therefore

$$\exp(D_u) - 1 = D_u \left(1 + D_u/2! + \dots + D_u^{k-1}/k! \right),$$

and the parenthesized operator is invertible. It follows that the translation stabilizer of F is the linear subspace

$$\text{Stab}(F) = \{u : F(X + u) = F(X)\} = \{u : D_u F = 0\}. \quad (\text{D1.2})$$

If (D1.1) has one solution, all its solutions form a coset of this subspace. Its codimension is

$$\dim_{\mathbb{F}_p} \text{span}\{\partial_1 F, \dots, \partial_k F\}. \quad (\text{D1.3})$$

If the dimension in (D1.3) is one, write $\partial_i F = a_i G$ for a nonzero polynomial G . When $a_i \neq 0$, multilinearity implies that G is independent of X_i . When $a_i = 0$, F is independent of X_i , and hence so is G . Thus G is constant and F is affine. In particular, every nonaffine F has a stabilizer of codimension at least two. For each pair of nonaffine slice patterns, (D1.1) consequently has at most p^{k-2} solutions.

It remains to determine the hypersurfaces supplied by affine F . Write $\pi(\epsilon) = (-1)^{\sum_i \epsilon_i}$ and $\eta(\epsilon) = \pi(\epsilon)v(\epsilon)$. The polynomial F_η is affine if and only if the unique multilinear representation of v is affine: direct expansion gives $F_\eta(X) = (-1)^k \sum_T a_T \prod_{i \in T} (X_i + 1)$ when $v(\epsilon) = \sum_T a_T \prod_{i \in T} \epsilon_i$, and substitution $X_i \mapsto X_i + 1$ preserves total degree. More explicitly, if

$$v(\epsilon) = c + \sum_i b_i \epsilon_i,$$

then

$$F_\eta(X) = (-1)^k \left(c + \sum_i b_i (X_i + 1) \right). \quad (\text{D1.4})$$

For $p \geq 5$, an affine Boolean function taking values only in $\{-1, 0, 1\}$ has one of the following forms:

- a constant (three choices);

- one nonzero slope $b_i = \pm 1$, with two admissible intercepts;
- one nonzero slope $b_i = \pm 2$, with one admissible intercept;
- two nonzero slopes $b_i, b_j \in \{-1, 1\}$, with one admissible intercept.

Indeed, after fixing a nonzero slope a , the values on the other coordinates must lie in $S \cap (S - a)$, where $S = \{-1, 0, 1\}$. This set has at most two elements, and has two only when $a = \pm 1$. An affine Boolean function with two values can depend on at most one coordinate in odd characteristic: otherwise the four values $x, x + b, x + c, x + b + c$, with $b, c \neq 0$, force $b = c$ and $2b = 0$. The displayed list and its intercept counts follow. Thus the number of affine slice patterns is

$$f_k = 3 + 6k + 4 \binom{k}{2} = 2k^2 + 4k + 3. \quad (\text{D1.5})$$

In (D1.1), replace η_2 by its negative, which is another allowed pattern. Formula (D1.4) shows that a successful pair must have the same slopes b_i and satisfy $\sum_i b_i t_i = c_2 - c_1$. Constant polynomials give only multiples of the universal certificate. The remaining cases give exactly the following distinct affine hyperplanes:

$$t_i = 0, 1, -1 \quad (1 \leq i \leq k), \quad t_i - t_j = 0, \quad t_i + t_j = 0 \quad (i < j). \quad (\text{D1.6})$$

Each equation is realized by allowed patterns. A vector outside all these hyperplanes has nonzero coordinates, none equal to ± 1 , and its coordinates belong to distinct unordered sign pairs. There are $p - 3$ choices for t_1 , $p - 5$ for t_2 , and in general $p - (2i + 1)$ for t_i . Hence the number outside (D1.6) is exactly

$$\prod_{i=1}^k (p - 2i - 1). \quad (\text{D1.7})$$

There are $3^{2^k} - f_k$ nonaffine slice patterns. Translation preserves degree, so any successful pair involving a nonaffine slice has two nonaffine slices. Simultaneously negating both slices gives the same solution set and has no fixed point on these pairs. There are therefore at most R_k such solution sets. Applying the codimension-two bound to each proves the asserted finite upper bound; retaining only (D1.6) and the zero-side loci proves the lower bound. Finally, the linear coefficient of the complement probability is

$$k + \sum_{i=1}^k (2i + 1) = k^2 + 3k.$$

For $k = 1$ every polynomial F_η is affine, so $R_1 = 0$. □

[write] 6 October 2026; sources 21–27. *Answers source 18's question on $d = 1$ in its leading coefficient* (item II.38); source 18 set the case aside after its Proposition II.66.3. Source 27's Proposition II.102.1 is the case $k = 1$, proved independently. The intake computed $\delta_{k,1}(p)$ from the full definition and found it equal to $L_k(p)$ for $k = 2$ at $p = 7, 11, 13, 17, 19, 23$ and for $k = 3$ at $p = 11, 13$, beyond this source's checks; that is evidence for its question on exactness (Research question V.92.2 here), not a proof.

[write] 7 October 2026; sources 47–53. *Sharpened by source 49* (Section II.220): rank-two certificates add no locus outside the signed-coordinate arrangement, so $0 \leq \delta_{k,1}(p) - L_k(p) \leq (1 - p^{-1})^k M_k^2 / (2p^3)$ and the coefficient of p^{-2} is $-k(3k^3 + 14k^2 + 15k - 14)/6$ (its Theorem II.222.4); $\delta_{k,1} = L_k$ exactly for $k \leq 3$ at every odd prime (Corollary II.222.5, Theorem II.223.1), which answers this source's exactness question for $k \leq 3$ (item II.42).

Connection with the balanced degree-one count. The Laurent-polynomial calculation has a particularly simple $q = 2$ specialization:

$$\text{CT } F_k^2 = 2k^2 + 8k + 9, \quad \text{CT } B_k^2 = 2k + 4.$$

Thus the number of balanced classes of degree at most one is $(k+1)^2$. Removing the one nonconstant degree-zero class gives $k^2 + 2k = 3k + 2\binom{k}{2}$ classes of degree exactly one. In the present two-cube setting these classes give precisely the distinct hypersurfaces in (D1.6). The k zero-side hypersurfaces account for the remaining k in the coefficient $k^2 + 3k$.

Scope. The theorem determines the optimal leading coefficient for the two-cube case omitted from the earlier $d \geq 2$ argument. It does not assert that all higher-codimension degeneracy is contained in (D1.6), and its explicit remainder is only a union bound.

II.86 Arrangements and their moment relations

Source 25, delivered Section 2.

Put $m = 2d$ and

$$w = (\underbrace{1, \dots, 1}_d, \underbrace{-1, \dots, -1}_d), \quad \mathcal{W} = w^\perp \subseteq \mathbb{F}_q^m.$$

An arrangement is specified by

$$h = (h_1, \dots, h_k) \in \mathbb{F}_q^k, \quad x = (x_{ji})_{1 \leq j \leq m, 1 \leq i \leq k}, \quad r \in \mathcal{W},$$

and its vertices are

$$\rho(\varepsilon, j) = (x_{j1} + \varepsilon_1 h_1, \dots, x_{jk} + \varepsilon_k h_k, r_j), \quad \varepsilon \in \{-1, 1\}^k. \quad (\text{II.86.1})$$

The relation $r \in \mathcal{W}$ is $r_1 + \dots + r_d = r_{d+1} + \dots + r_{2d}$. There are $q^k q^{2dk} q^{2d-1} = q^D$ arrangements. The usual vertices $y_{ji} + e_i H_i$, $e_i \in \{0, 1\}$, are obtained by $x_{ji} = y_{ji} + H_i/2$ and $h_i = H_i/2$. This is a bijection of parameter spaces, not an approximation or a restriction.

Write $[k] = \{1, \dots, k\}$ and $\pi(\varepsilon) = \prod_{i=1}^k \varepsilon_i$. The distinguished coefficient pattern is

$$\eta_0(\varepsilon, j) = w_j \pi(\varepsilon). \quad (\text{II.86.2})$$

The Boolean convention $(-1)^{\sum e_i}$ differs by the harmless global factor $(-1)^k$.

For a ternary pattern $\eta : \{-1, 1\}^k \times [m] \rightarrow \{-1, 0, 1\} \subset \mathbb{F}_q$ define

$$M_A(\eta) = \sum_{j=1}^m \sum_{\varepsilon} \eta(\varepsilon, j) \prod_{i \in A} (x_{ji} + \varepsilon_i h_i), \quad (\text{II.86.3})$$

$$M_{A,r}(\eta) = \sum_{j=1}^m r_j \sum_{\varepsilon} \eta(\varepsilon, j) \prod_{i \in A} (x_{ji} + \varepsilon_i h_i), \quad A \subseteq [k]. \quad (\text{II.86.4})$$

These are exactly all coordinate-monomial moments in $k+1$ variables. In particular, $A = \emptyset$ is included.

Definition II.86.1. An arrangement is *degenerate* if some ternary η outside the one-dimensional space $\mathbb{F}_q\eta_0$ satisfies $M_A(\eta) = M_{A,r}(\eta) = 0$ for every $A \subseteq [k]$. Such an η is a witness. A pattern is *block-parity* if $\eta(\varepsilon, j) = a_j\pi(\varepsilon)$ for every j .

Orthogonality on the sign cube shows that the block moment of π is zero unless $A = [k]$, and is then $2^k \prod_i h_i$. Thus η_0 annihilates all moments because both $\sum_j w_j = 0$ and $\sum_j w_j r_j = 0$. Since the values of η_0 are ± 1 , its only ternary scalar multiples are $0, \eta_0, -\eta_0$.

Lemma II.86.2 (Coincidences and zero sides). *Every arrangement with a zero sidelength is degenerate. Every nondegenerate arrangement has distinct labelled vertices.*

Proof. If two labelled vertices coincide, put coefficient 1 at one and -1 at the other, and zero elsewhere. Every monomial moment cancels. This pattern cannot be a nonzero scalar multiple of η_0 , whose support consists of all $n \geq 8$ labels. If $h_i = 0$, the vertices within one block whose signs differ only at i coincide, so the first assertion is a special case of the second argument. $\square$

A function $\phi : B \rightarrow \mathbb{F}_q$ respects an arrangement in B if

$$\sum_{j, \varepsilon} w_j \pi(\varepsilon) \phi(\rho(\varepsilon, j)) = 0. \quad (\text{II.86.5})$$

We will need an elementary upper bound for the total number of arrangements contained in an arbitrary set.

Lemma II.86.3 (Density cap). *If $B \subseteq \mathbb{F}_q^{k+1}$ has size βq^{k+1} , then it contains at most $\beta^{2d-1} q^D$ labelled d -arrangements.*

Proof. Choose h , the all-minus vertex of each of the first $2d-1$ blocks, and the k spatial coordinates of the all-minus vertex of the last block. The cross-section relation determines the remaining coordinate of the last block. These choices determine the entire arrangement. For an arrangement in B , the first $2d-1$ selected vertices must be in B . There are at most $q^k |B|^{2d-1} q^k = \beta^{2d-1} q^D$ choices. Ignoring the other containment requirements only increases this upper bound. $\square$

II.87 Minimal Walsh layers give independent linear equations

Source 25, delivered Section 3.

For $S \subseteq [k]$, let $\chi_S(\varepsilon) = \prod_{i \in S} \varepsilon_i$ and put

$$c_{j,S} = 2^{-k} \sum_{\varepsilon} \eta(\varepsilon, j) \chi_S(\varepsilon). \quad (\text{II.87.1})$$

Here 2^{-k} denotes inversion in the field. Orthogonality gives

$$\sum_{\varepsilon} \chi_S(\varepsilon) \chi_T(\varepsilon) = \begin{cases} 2^k & S = T, \\ 0 & S \neq T, \end{cases} \quad \eta(\varepsilon, j) = \sum_{S \subseteq [k]} c_{j,S} \chi_S(\varepsilon). \quad (\text{II.87.2})$$

Indeed, a coordinate in $S \triangle T$ makes the sum cancel in pairs; otherwise every summand is one. This proves inversion as well.

Use $h_S = \prod_{i \in S} h_i$ and $x_{j,S} = \prod_{i \in S} x_{ji}$. Expanding the product in a block moment and applying (II.87.1) gives

$$\sum_{\varepsilon} \eta(\varepsilon, j) \prod_{i \in A} (x_{ji} + \varepsilon_i h_i) = 2^k \sum_{S \subseteq A} c_{j,S} h_S x_{j,A \setminus S}. \quad (\text{II.87.3})$$

This identity is valid for arbitrary field-valued patterns, not only ternary ones.

Lemma II.87.1 (Minimal-layer linearization). *Suppose $h_i \neq 0$ for every i . For a nonzero pattern let*

$$s = \min\{|S| : c_{j,S} \neq 0 \text{ for some } j\}.$$

If $s < k$, then the equations $M_A(\eta) = 0$ with $|A| = s + 1$ include $k - s$ independent affine-linear equations in the $2dk$ base coordinates. For every fixed r , the probability over uniform x that all moments vanish is therefore at most $q^{-(k-s)}$.

Proof. All coefficients of degree below s vanish, so (II.87.3), for $|A| = s + 1$, contains only terms linear in x and constant terms. Choose j_0, S_0 with $|S_0| = s$ and $c_{j_0, S_0} \neq 0$. For each $i \notin S_0$, use the equation indexed by $A_i = S_0 \cup \{i\}$, and use $x_{j_0 i}$ as its pivot variable. The coefficient of this variable is $2^k c_{j_0, S_0} h_{S_0} \neq 0$. If $i' \notin S_0$ and $i' \neq i$, then $i' \notin A_i$, so $x_{j_0 i'}$ does not occur in that equation. The indicated square minor is diagonal and nonsingular.

After every other base coordinate has been fixed, these $k - s$ equations determine the pivot coordinates uniquely. They therefore have exactly $q^{2dk - (k-s)}$ simultaneous solutions; imposing the remaining moment equations can only reduce this number. No polynomial zero-counting estimate is needed. $\square$

Remark II.87.2. Gowers's Lemma 15.3 is sharp for a single arbitrary multilinear polynomial: a product of k coordinates can vanish on approximately k/q of the space. Lemma II.87.1 does not contradict that fact. It exploits the simultaneous moment system, in which a lowest nonzero Walsh layer provides specially chosen *linear* equations. It strengthens the use of Lemma 15.3 inside the arrangement argument, not the original fibre bound without extra hypotheses.

II.88 Non-parity witnesses have codimension at least two

Source 25, delivered Section 4.

Fix nonzero h and a ternary pattern that is not block-parity. Its minimal layer has $s \leq k - 1$. If $s \leq k - 2$, Lemma II.87.1 already gives a q^{-2} bound, uniformly in the cross-sections. Only $s = k - 1$ needs a second argument.

In this remaining case write

$$\eta(\varepsilon, j) = a_j \pi(\varepsilon) + \sum_{i=1}^k b_{ji} \pi(\varepsilon) \varepsilon_i, \quad a_j = c_{j, [k]}, \quad b_{ji} = c_{j, [k] \setminus \{i\}}. \quad (\text{II.88.1})$$

At least one vector $b_i = (b_{1i}, \dots, b_{mi})$ is nonzero. The coefficients a_j, b_{ji} are field elements; they need not be ternary. The moments indexed by $[k] \setminus \{i\}$ imply

$$\sum_j b_{ji} = 0, \quad b_i \cdot r = 0 \quad (1 \leq i \leq k). \quad (\text{II.88.2})$$

After division by $2^k h_{[k]}$, the two top moments become

$$\sum_{j,i} b_{ji} \frac{x_{ji}}{h_i} + \sum_j a_j = 0, \quad (\text{II.88.3})$$

$$\sum_{j,i} r_j b_{ji} \frac{x_{ji}}{h_i} + \sum_j r_j a_j = 0. \quad (\text{II.88.4})$$

These equations also display why keeping moments containing the final coordinate is useful.

Theorem II.88.1 (Sharp single-witness bound). *Let $d \geq 2$, $k \geq 1$, and let the field have odd characteristic. For each fixed nonzero sidelength vector and each non-block-parity ternary pattern, uniform $r \in \mathcal{W}$ and uniform x satisfy all moment equations with probability at most*

$$e_d(q) = q^{-2} + (q-1)q^{-2d}. \quad (\text{II.88.5})$$

This bound is attained, for every k , by

$$\eta(\varepsilon, j) = w_j \pi(\varepsilon) \varepsilon_{i_0} \quad (\text{II.88.6})$$

for any fixed $i_0 \in [k]$.

Proof. The case $s \leq k-2$ follows from Lemma II.87.1. Suppose $s = k-1$. If a constant constraint in (II.88.2) fails, there are no solutions. Otherwise put

$$e = \dim \text{span}(w, b_1, \dots, b_k) - 1.$$

When $e \geq 1$, the equations $b_i \cdot r = 0$ cut out a subspace of codimension e in $\mathcal{W}$. Their probability is q^{-e} . For every admissible r , equation (II.88.3) is a nonconstant affine-linear equation in x , of probability q^{-1} . Thus the simultaneous probability is at most $q^{-e-1} \leq q^{-2}$.

When $e = 0$, write $b_i = t_i w$, where some $t_i \neq 0$. The cross-section constraints are automatic. The coefficient rows of (II.88.3) and (II.88.4) are linearly dependent exactly when all r_j are equal: for an i with $t_i \neq 0$, comparison of the coefficients of x_{ji} forces r_j to be the same scalar for every j . Outside that one-dimensional diagonal subspace of $\mathcal{W}$, the two equations have independent rows and hence exactly q^{2dk-2} solutions for any right hand sides. On the diagonal they have at most q^{2dk-1} solutions. The diagonal occupies a proportion $q^{-(2d-2)}$ of $\mathcal{W}$. Therefore the probability is at most

$$(1 - q^{-(2d-2)})q^{-2} + q^{-(2d-2)}q^{-1} = q^{-2} + (q-1)q^{-2d}.$$

For (II.88.6), we have $a_j = 0$, $b_{i_0} = w$, and $b_i = 0$ for $i \neq i_0$. All nontrivial conditions are exactly

$$\sum_j w_j x_{ji_0} = 0, \quad \sum_j w_j r_j x_{ji_0} = 0.$$

They have rank two off the diagonal and rank one on it, and they are always consistent. Thus equality holds in the displayed count. $\square$

[write] **6 October 2026; sources 21–27.** *Credit.* First printed as source 18's Proposition II.66.3 (with the same witness up to the Boolean–sign convention); also source 22's Lemma II.82.3 at degree one and source 27's Proposition II.94.3 (every characteristic).

Example II.88.2 (The restriction $d \geq 2$ is necessary). For $d = 1$, the cross-section relation forces $r_1 = r_2$. The same witness (II.88.6) gives just $x_{1i_0} = x_{2i_0}$, a probability q^{-1} event. Consequently a blanket codimension-two assertion for non-parity witnesses is false for $d = 1$. This is a structural obstruction, not a weakness of the proof. The formula in (II.88.5) itself becomes q^{-1} in this case.

For the final count it is useful to exclude a set of cross-sections that will already be counted as degenerate.

Lemma II.88.3 (Removing the diagonal exception). *Let $\mathcal{R} \subseteq \mathcal{W}$ omit every constant vector. Under the hypotheses of Theorem II.88.1, for each fixed non-parity pattern and each fixed nonzero h ,*

$$\mathbb{P}_{r \in \mathcal{W}, x}(r \in \mathcal{R} \text{ and all moments vanish}) \leq q^{-2}. \quad (\text{II.88.7})$$

For the witness (II.88.6), equality holds with the right side replaced by $(|\mathcal{R}|/|\mathcal{W}|)q^{-2}$.

Proof. For $s \leq k - 2$ and for $e \geq 1$, the corresponding bounds in the previous proof were already q^{-2} . For $e = 0$, the only rank-one exception was the diagonal, now excluded. The final assertion follows from the two independent homogeneous equations for the sharp witness. $\square$

Notice the normalization: (II.88.7) is a *joint* probability with r uniform in all of $\mathcal{W}$. Conditioning r to belong to $\mathcal{R}$ requires division by $|\mathcal{R}|/|\mathcal{W}|$.

II.89 The cross-section hyperplanes and their exact number

Source 25, delivered Section 5.

Suppose all sidelengths are nonzero and η is block-parity, $\eta(\varepsilon, j) = a_j \pi(\varepsilon)$. Since the pattern is ternary, every a_j belongs to $\{-1, 0, 1\}$. Its moment conditions reduce exactly to

$$\sum_{j=1}^m a_j = 0, \quad a \cdot r = 0. \quad (\text{II.89.1})$$

The first equality follows from the top spatial moment, and the second from the top moment including the last coordinate. All other moments vanish identically. The condition $p > 2d$ makes the first equality in the field equivalent to the integer equality: its integer sum has absolute value at most $2d$.

Exclude $a = 0, w, -w$, which give the distinguished scalar multiples. Every remaining balanced ternary vector defines a hyperplane

$$H_a = \{r \in \mathcal{W} : a \cdot r = 0\}. \quad (\text{II.89.2})$$

Let $\mathcal{H}_d$ be the set of *distinct* such hyperplanes. Each contains the diagonal of constant cross-sections, since $\sum a_j = 0$.

Theorem II.89.1 (Projective enumeration). *For $d \geq 1$ and $p > \max(3, 2d)$,*

$$|\mathcal{H}_d| = H_d = \frac{T_{2d} - 2\binom{2d}{d} + 1}{2}.$$

For $d \geq 2$, this family is nonempty.

Proof. Two nonzero restrictions $a|_{\mathcal{W}}$ and $a'|_{\mathcal{W}}$ have the same kernel exactly when

$$a' = \lambda a + \mu w, \quad \lambda \neq 0. \quad (\text{II.89.3})$$

This is the elementary fact that two nonzero linear forms with the same kernel are proportional, together with $\mathcal{W}^\perp = \mathbb{F}_q w$. Set $z_j = w_j a_j$. Then $z \in \{-1, 0, 1\}^{2d}$, $\sum_j w_j z_j = 0$, and (II.89.3) becomes

$$z' = \lambda z + \mu \mathbf{1}. \quad (\text{II.89.4})$$

We must count nonconstant balanced ternary vectors modulo affine changes of their values.

There are T_{2d} balanced vectors: multiplying their coordinates by w_j bijects them with the integer-balanced vectors a in (II.89.1). Exactly three are constant, and they correspond to $a = 0, w, -w$.

First suppose z uses exactly two values $u < v$. If S is the set of coordinates at the larger value, then

$$\sum_j w_j z_j = (v - u) \sum_{j \in S} w_j.$$

It vanishes exactly when S contains equally many coordinates from the first and second halves. The number of nonempty proper such subsets is

$$B_d = \sum_{t=0}^d \binom{d}{t}^2 - 2 = \binom{2d}{d} - 2. \quad (\text{II.89.5})$$

The binomial identity follows by counting d -element subsets of two disjoint d -element sets, according to the number chosen from the first. There are three choices of a two-element value set, hence $3B_d$ vectors of this type. Each affine class consists of six vectors: three choices of value set and the two choices of which part gets the larger value. Thus there are $B_d/2$ two-value classes.

If z uses all three values, an affine change preserving ternarity must permute $\{-1, 0, 1\}$. Summing the three images gives $3\mu = 0$; as $p > 3$, we have $\mu = 0$. The image of 1 now forces $\lambda = 1$ or -1 . Consequently every three-value class consists exactly of z and $-z$. No invertible affine change can alter the number of distinct values. The three-value classes therefore number $(T_{2d} - 3 - 3B_d)/2$. Adding the two contributions gives

$$\frac{B_d}{2} + \frac{T_{2d} - 3 - 3B_d}{2} = \frac{T_{2d} - 2\binom{2d}{d} + 1}{2}.$$

For $d \geq 2$, for example $a = (1, -1, 0, \dots, 0)$ is balanced and not a multiple of w , so at least one hyperplane occurs. $\square$

[write] **6 October 2026; sources 21–27.** *Credit.* H_d is source 18's C_d (Lemma II.66.6), source 22's C_d and source 27's L_d . The intake counted the distinct kernels by rank tests and found 0, 4, 51 for $d = 1, 2, 3$ over $\mathbb{F}_{101}$, and 57 for $d = 3$ over $\mathbb{F}_5$, as Example II.89.2 and source 27's Example II.96.4 say.

II.89.1 An explicit canonical representative

Eliminate the last cross-section coordinate using $r_m = \sum_{i < m} w_i r_i$. The restricted normal of H_a is

$$v(a) = (a_i + a_m w_i)_{i < m}. \quad (\text{II.89.6})$$

Over the rationals one divides this nonzero integer vector by the gcd of its entries and chooses its first nonzero entry positive. Over a prime field one divides by its first nonzero entry. These are the two independent representations used in the supplied enumeration check. The closed formula does not rely on either computation.

d	T_{2d}	$\binom{2d}{d}$	H_d
1	3	2	0
2	19	6	4
3	141	20	51
4	1 107	70	484
5	8 953	252	4 225
6	73 789	924	35 971
7	616 227	3 432	304 682
8	5 196 627	12 870	2 585 444

Table II.89.1: The exact number of distinct parity hyperplanes. Rows through $d = 5$ were also checked by direct projective enumeration; the remaining rows are evaluations of the proved formula.

Example II.89.2 (Why a large field cardinality is not enough). For $d = 3$ in characteristic 5, the integer sums ± 5 also vanish in the field. They arise from five coefficients equal to 1 and one zero, or from their negatives. There are twelve such vectors. Each gives a three-value $z = w \cdot a$, so the preceding affine classification groups them only in opposite pairs; they cannot coincide with integer-balanced classes. There are therefore six extra hyperplanes:

$$|\mathcal{H}_3(\mathbb{F}_5)| = 57 \neq 51.$$

The same phenomenon persists in every extension field of characteristic 5. The hypothesis is on the characteristic, not merely on q .

II.90 The global count and the geometry of the exceptional set

Source 25, delivered Section 6.

Let

$$\mathcal{V}_d = \bigcup_{H \in \mathcal{H}_d} H \subseteq \mathcal{W}, \quad v_d(q) = \frac{|\mathcal{V}_d|}{q^{2d-1}}.$$

Call an arrangement *base-exceptional* if some $h_i = 0$ or if $r \in \mathcal{V}_d$. The word “base” here refers to the explicitly classified exception set, not to the spatial base coordinates x . The probability of this event is exactly

$$b_{k,d}(q) = 1 - (1 - q^{-1})^k (1 - v_d(q)). \quad (\text{II.90.1})$$

Every base-exceptional arrangement is degenerate. For zero sides this is Lemma II.86.2; for the cross-section part, use the corresponding block-parity witness, whose moment identities hold for all spatial bases and all sidelengths.

Theorem II.90.1 (Residual count after all hypersurfaces are removed). *Under (F.30.2),*

$$b_{k,d}(q) + (1 - b_{k,d}(q))q^{-2} \leq \frac{Z_{k,d}(q)}{q^D} \leq b_{k,d}(q) + (1 - q^{-1})^k U_{k,d} q^{-2}. \quad (\text{II.90.2})$$

In particular, the residual degeneracy outside the classified hyperplanes is genuinely of order q^{-2} for fixed k, d .

Proof. Outside the base-exceptional set, all sides are nonzero and no block-parity witness is available. There are 3^n ternary patterns in all, and precisely 3^{2d} block-parity patterns. The remaining patterns come in opposite pairs defining identical vanishing events. Thus there are $U_{k,d}$ sign classes to consider.

The regular cross-section set $\mathcal{W} \setminus \mathcal{V}_d$ omits the constant vectors, since $\mathcal{H}_d$ is nonempty and all its hyperplanes contain them. For each non-parity sign class, Lemma II.88.3 bounds the joint probability of vanishing and regular cross-sections by q^{-2} for every nonzero h . Summing these bounds and multiplying by $(1 - q^{-1})^k$ proves the upper bound. This is a union bound, so overlaps cause no problem.

For the lower bound fix one witness (II.88.6). At every nonzero h and every regular r , its two equations have rank two and have exactly probability q^{-2} over x . All these configurations are degenerate and none is base-exceptional. Their probability is $(1 - b_{k,d}(q))q^{-2}$, proving the lower bound.

Finally $v_d(q) \leq H_d/q$, so $b_{k,d}(q) = O_{k,d}(q^{-1})$. The extra term on the lower side is consequently $q^{-2} + O_{k,d}(q^{-3})$, whereas the extra term on the upper side is $O_{k,d}(q^{-2})$. $\square$

Proof of the counting part of Theorem F.30.1. The base-exceptional set is the union of $L_{k,d} = k + H_d$ distinct linear hyperplanes in the D -dimensional parameter space $\mathbb{F}_q^k \times \mathbb{F}_q^{2dk} \times \mathcal{W}$. Each hyperplane has probability q^{-1} , and each pair of distinct hyperplanes has intersection probability q^{-2} . The union bound and the first two terms of inclusion–exclusion give

$$\frac{L_{k,d}}{q} - \frac{\binom{L_{k,d}}{2}}{q^2} \leq b_{k,d}(q) \leq \frac{L_{k,d}}{q}. \quad (\text{II.90.3})$$

The lower estimate may be negative for small q , which does not affect its validity. Combine (II.90.3) with Theorem II.90.1, and discard $(1 - q^{-1})^k \leq 1$ on the upper side. Multiplication by q^D proves (F.30.5); division by q^{D-1} proves (F.30.6). $\square$

II.90.1 Codimension-one components over an infinite field

The codimension assertion is stronger than a finite counting heuristic, so we give its geometric justification. Work over an algebraically closed field K of characteristic zero or greater than $2d$. The ambient space is affine D -space after eliminating r_m . For each of the finitely many ternary patterns, its moment-vanishing locus is a closed algebraic set. The degeneracy locus is the finite union of these sets, with the three distinguished patterns excluded.

The zero-side hyperplanes and parity hyperplanes lie in this locus. On the open set $h_1 \cdots h_k \neq 0$, a non-parity witness with minimal layer $s \leq k - 2$ has at least two independent affine equations with a nonsingular pivot minor. Its dimension is at most $D - 2$. For $s = k - 1$ and $e \geq 1$, the cross-section equations give codimension at least one and (II.88.3) gives another independent condition in spatial variables. For $e = 0$, the generic cross-section fibre has two independent spatial equations; the exceptional diagonal has codimension $2d - 2$ in $\mathcal{W}$ and still has one nonconstant spatial equation. Its total codimension is at least $2d - 1 \geq 3$.

Thus every non-parity witness has codimension at least two away from the zero-side hyperplanes. Taking a finite union cannot create a new codimension-one component. Conversely, each of the listed distinct hyperplanes is irreducible of codimension one and is already contained in the degeneracy locus. This proves the remaining assertion of Theorem F.30.1 without deriving geometry from point counts.

II.90.2 The first nontrivial example

For $d = 2$, on $r_1 + r_2 = r_3 + r_4$ the four hyperplanes can be written

$$r_1 = r_2, \quad r_3 = r_4, \quad r_1 = r_3, \quad r_1 = r_4. \quad (\text{II.90.4})$$

Other balanced ternary equations give one of these or the defining relation. The four hyperplanes share the diagonal line; in the quotient by that line, they are four distinct lines in a two-dimensional space. Their union has size $4(q^2 - q) + q = 4q^2 - 3q$, so

$$v_2(q) = \frac{4}{q} - \frac{3}{q^2}, \quad b_{k,2}(q) = 1 - (1 - q^{-1})^k(1 - 4q^{-1} + 3q^{-2}). \quad (\text{II.90.5})$$

For example, $b_{1,2}(q) = 5q^{-1} - 7q^{-2} + 3q^{-3}$. This is an exact count of the classified hyperplanes, not an assertion that all degeneracies have been counted exactly.

II.91 Explicit ambient-size savings

Source 25, delivered Section 7.

Suppose a later argument can tolerate at most a fraction $\tau > 0$ of ambient arrangements being degenerate. Gowers's bound (F.30.1) gives the sufficient condition $q \geq k3^n/\tau$. The refined count instead gives the following.

Corollary II.91.1 (Two-scale sufficient cutoff). *Subject to (F.30.2), the condition*

$$q \geq \max \left\{ \frac{2L_{k,d}}{\tau}, \sqrt{\frac{2U_{k,d}}{\tau}} \right\} \quad (\text{II.91.1})$$

implies $Z_{k,d}(q) \leq \tau q^D$.

Proof. The first requirement makes $L_{k,d}/q \leq \tau/2$ and the second makes $U_{k,d}/q^2 \leq \tau/2$. Apply Theorem F.30.1. $\square$

The large ternary count now enters only through its square root, while the first-order contribution uses the much smaller exact coefficient. For extremely small τ , the first term in (II.91.1) can dominate; thus it would be inaccurate to describe the entire threshold as always a square-root improvement.

k	$n = 2^{k+4}$	$\log_{10}(k3^n/\tau)$	$\log_{10}$ of (II.91.1)	controlling term
1	32	21.268	12.714	first-order
2	64	36.837	18.268	second-order
3	128	67.549	33.536	second-order
4	256	128.745	64.072	second-order

Table II.91.1: Illustrative sufficient thresholds for $d = 8$ and $\tau = 10^{-6}$, before rounding to a permissible field cardinality. These are comparisons of proven upper-bound criteria, not measured minimal field sizes. Values are rounded to three decimals.

The tolerated fraction τ in the actual Riesz-product application can be far smaller than 10^{-6} , and the resulting sufficient field-size criterion far larger than this illustration. The next section supplies its precise value instead of silently treating “sufficiently large” as a parameter-free condition.

II.92 A finite-field random-selection theorem with exact parameters

Source 25, delivered Section 8.

We now generalize the selection step to arbitrary $d \geq 2$ and finite fields satisfying (F.30.2). The construction is the Riesz-product idea of Gowers [1, Section 15]; the proof below keeps exact survival probabilities and handles repeated labels explicitly.

Set

$$a_n = 2^{-n}, \quad b_n = 2^{-n}(1 + 2^{1-n}), \quad R_n = b_n/a_n = 1 + 2^{1-n}. \quad (\text{II.92.1})$$

Let $0 < \alpha, \beta, \eta \leq 1$. Define the integer

$$t = \min\{u \in \mathbb{Z}_{\geq 0} : R_n^u \geq 4/(\alpha\eta)\} = \left\lceil \frac{\log(4/(\alpha\eta))}{\log R_n} \right\rceil, \quad (\text{II.92.2})$$

and let

$$\tau = \frac{\alpha\eta}{4} \beta^{2d-1} b_n^t. \quad (\text{II.92.3})$$

For numerical implementation the minimum characterization avoids any rounding ambiguity at an integer logarithmic boundary.

Theorem II.92.1 (Explicit arrangement purification). *Let $B \subseteq \mathbb{F}_q^{k+1}$ have size βq^{k+1} , and let $\phi : B \rightarrow \mathbb{F}_q$ respect at least $\alpha\beta^{2d-1}q^D$ arrangements in B . Suppose (F.30.2) holds and*

$$\frac{L_{k,d}}{q} + \frac{U_{k,d}}{q^2} \leq \tau. \quad (\text{II.92.4})$$

Then a subset $B' \subseteq B$ exists which contains at least

$$\frac{\alpha}{2} \beta^{2d-1} b_n^t q^D \quad (\text{II.92.5})$$

respected arrangements, and the number of its non-respected arrangements is at most η times its number of respected arrangements. In particular, the proportion respected is at least $1/(1+\eta) \geq 1-\eta$. Condition (II.91.1) with (II.92.3) is a sufficient explicit condition for (II.92.4).

[write] **6 October 2026; sources 21–27.** *Against source 18 and the formal statement (checked at the write). Source 18's Corollary II.66.9 treats $d = 8$ over prime fields with the retention $(\alpha\eta/4)^{E_k} \beta^{15}$ of `lemma_15_5` (`Sections14_15.lean`, line 414; `arrangementSelectionExponent`, line 409, is $E_k = 2^{2^{k+4}+k+3}$); this theorem treats every $d \geq 2$ and every finite field of characteristic $p > 2d$ and keeps the exact number t of factors. At $d = 8$, $n = 2^{k+4}$ and $E_k = n2^{n-1}$, so (II.92.12) says $K_n/E_k \rightarrow \log 2$: the power-form retention $b_n(\alpha\eta/4)^{K_n}$ improves the leading factor of the exponent, not its form, at the cost of the prefactor $(\alpha/2)b_n$. Numerically $K_n/E_k = 0.69315$ already for $k = 1, 2, 3$. Not formalized; `lemma_15_5_holds` (`Proofs15Restriction.lean`, line 3086) proves the formal statement.*

II.92.1 One Riesz factor

Choose a nontrivial additive character $\psi : \mathbb{F}_q \rightarrow \mathbb{C}$. For example, choose a nonzero $\mathbb{F}_p$ -linear functional $\ell : \mathbb{F}_q \rightarrow \mathbb{F}_p$ and take $\psi(z) = \exp(2\pi i \ell(z)/p)$, where the exponential is independent of the integer representative of $\ell(z)$. For independent uniform u and s_A in $\mathbb{F}_q$, with $A \subseteq [k+1]$, put

$$P(z) = \frac{1 + \Re \psi(u\phi(z) + \sum_{A \subseteq [k+1]} s_A z_A)}{2}, \quad z_A = \prod_{i \in A} z_i. \quad (\text{II.92.6})$$

Thus $0 \leq P(z) \leq 1$. The empty monomial is included. The expansion of this factor is

$$P(z) = \frac{1}{2} + \frac{1}{4}\psi(u\phi(z) + \sum_A s_A z_A) + \frac{1}{4}\psi(-u\phi(z) - \sum_A s_A z_A). \quad (\text{II.92.7})$$

All coefficients are nonnegative.

Orthogonality says that $\mathbb{E}_s \psi(sc)$ is 1 if $c = 0$ and 0 otherwise. For $c \neq 0$, multiplication by c permutes the field, and the average of a nontrivial character is zero: translating its sum by an element on which it is not one multiplies the sum by a scalar different from one without changing it. Applying this to every s_A , the surviving terms in the product of (II.92.7) over arrangement labels are exactly the ternary patterns whose monomial moments vanish. Averaging over u additionally requires $\sum_{j,\varepsilon} \eta(\varepsilon, j) \phi(\rho(\varepsilon, j)) = 0$.

For a nondegenerate arrangement, the only possibilities are 0, η_0 , $-\eta_0$. The zero pattern contributes 2^{-n} . Each of the two nonzero patterns has full support and contributes 4^{-n} precisely when ϕ respects the arrangement. Hence

$$\mathbb{E} \prod_{j,\varepsilon} P(\rho(\varepsilon, j)) = \begin{cases} a_n, & \phi \text{ does not respect the arrangement,} \\ b_n, & \phi \text{ respects the arrangement.} \end{cases} \quad (\text{II.92.8})$$

For a degenerate but respected arrangement the expectation is at least b_n : the same three terms still survive, and all additional surviving terms have nonnegative coefficients.

II.92.2 Several factors and repeated vertices

Take t independent choices of the parameters in (II.92.6), giving factors $P_1, \dots, P_t$. Conditional on these choices, put each *point* $z \in B$ in B' independently with probability

$$p(z) = \prod_{\ell=1}^t P_\ell(z).$$

A nondegenerate arrangement has distinct vertices by Lemma II.86.2. Its conditional survival probability is therefore the product of $p(z)$ over its labels. Independence between the t parameter choices and (II.92.8) give survival probability a_n^t or b_n^t according as the arrangement is not respected or respected.

For a repeated-vertex arrangement the conditional survival probability is a product over *distinct* points. Since all factors belong to $[0, 1]$, this is at least the product over all labelled occurrences. The preceding nonnegative Fourier expansion therefore proves that every respected arrangement, even a degenerate one, survives with probability at least b_n^t . It would be incorrect to assert equality of the two products when labels collide.

II.92.3 The positive-potential argument

Proof of Theorem II.92.1. Let $M = \beta^{2d-1} q^D$. Let X count respected arrangements in the selected set and Y count non-respected ones. By the preceding survival calculations, the hypothesis, and Lemma II.86.3,

$$\mathbb{E}X \geq \alpha M b_n^t, \quad \mathbb{E}Y \leq M a_n^t + Z_{k,d}(q). \quad (\text{II.92.9})$$

For the second inequality, treat nondegenerate arrangements using their exact probability a_n^t , and bound the total contribution of degenerate arrangements by their ambient number.

By the definition of t , $a_n^t \leq (\alpha\eta/4)b_n^t$. By Theorem F.30.1, (II.92.4), and (II.92.3), $Z_{k,d}(q) \leq (\alpha\eta/4)Mb_n^t$. Consequently

$$\mathbb{E}(X - \eta^{-1}Y) \geq \alpha Mb_n^t - \frac{\alpha}{4}Mb_n^t - \frac{\alpha}{4}Mb_n^t = \frac{\alpha}{2}Mb_n^t.$$

There is an outcome with $X - \eta^{-1}Y \geq (\alpha/2)Mb_n^t$. For that outcome, X is at least the right side of (II.92.5), and $Y \leq \eta X$. Thus $X/(X + Y) \geq 1/(1 + \eta) \geq 1 - \eta$. $\square$

II.92.4 A power-form bound and the remaining loss

Put

$$K_n = \frac{\log(1/b_n)}{\log R_n}. \quad (\text{II.92.10})$$

Since $t < 1 + \log(4/(\alpha\eta))/\log R_n$ and $0 < b_n < 1$,

$$b_n^t \geq b_n \left(\frac{\alpha\eta}{4} \right)^{K_n}. \quad (\text{II.92.11})$$

This gives a power-form lower bound for (II.92.5) with no unspecified constants. Direct expansion of $\log(1 + 2^{1-n})$ gives

$$K_n = (\log 2) n2^{n-1} + O(n) \quad (n \rightarrow \infty). \quad (\text{II.92.12})$$

For $d = 8$, the exponent recorded in the repository for the original coarse selection bound is $2^{2k+4+k+3} = n2^{n-1}$ [114]. Keeping the exact logarithm changes the leading numerical factor in this local exponent, but not its exponential dependence on n . The very small likelihood advantage $R_n - 1 = 2^{1-n}$ remains the main loss in this particular selection mechanism.

The theorem neither assumes an improved inverse theorem nor silently imports a result about prime fields into a composite ring. It is valid over finite *fields*; the additive character averages are taken over that field, and the rank calculations are field calculations.

II.93 Labelled arrangements and their moment relations

Source 27, delivered Section 2.

II.93.1 The parameter space

Let $s = (s_1, \dots, s_n)$ have its first d coordinates equal to 1 and its last d coordinates equal to -1 . All signs are interpreted in the field when the characteristic is two. Put

$$U = \{u = (u_1, \dots, u_n) \in \mathbb{F}_q^n : s \cdot u = 0\}. \quad (\text{II.93.1})$$

It has dimension $n - 1$ because $s \neq 0$. A labelled arrangement consists of

$$h \in \mathbb{F}_q^k, \quad x_1, \dots, x_n \in \mathbb{F}_q^k, \quad u \in U.$$

Its labelled vertices, indexed by $\varepsilon \in \{0, 1\}^k$ and $1 \leq j \leq n$, are

$$v_{\varepsilon,j} = (x_{j1} + \varepsilon_1 h_1, \dots, x_{jk} + \varepsilon_k h_k, u_j) \in \mathbb{F}_q^{k+1}. \quad (\text{II.93.2})$$

There are $q^k q^{nk} q^{n-1} = q^D$ parameter tuples, as asserted in (F.32.1). All probabilities below refer to the uniform measure on this parameter space. In particular the bases, side coordinates, and the uniform cross section in U are independent.

For a coefficient array $\eta = (\eta_{\varepsilon,j})$ and a subset $A \subseteq [k+1]$, define the squarefree moment

$$M_A(\eta) = \sum_{j=1}^n \sum_{\varepsilon \in \{0,1\}^k} \eta_{\varepsilon,j} \prod_{i \in A} v_{\varepsilon,j,i}. \quad (\text{II.93.3})$$

The empty product is 1; thus $M_{\emptyset}(\eta) = \sum \eta_{\varepsilon,j}$.

II.93.2 The intrinsic relation and the alphabet

Write

$$\pi(\varepsilon) = (-1)^{|\varepsilon|}, \quad e_{\varepsilon,j} = \pi(\varepsilon)s_j.$$

For $A \subsetneq [k]$, the alternating cube sum of the associated monomial vanishes. For $A = [k]$ it is $(-1)^k \prod_i h_i$, independent of the base. The same statement with the final coordinate present acquires the factor u_j . Since $\sum s_j = 0$ and $\sum s_j u_j = 0$, this proves

$$M_A(e) = 0 \quad (A \subseteq [k+1]) \quad (\text{II.93.4})$$

in every characteristic, including when some $h_i = 0$.

Let $D_p \subseteq \mathbb{F}_p \subseteq \mathbb{F}_q$ be the image of $\{-1, 0, 1\}$. In characteristic two this alphabet has two elements, not three.

Definition II.93.1. An arrangement is *degenerate* if there is an array $\eta \in D_p^{\{0,1\}^k \times [n]}$ such that

$$\eta \notin \mathbb{F}_q e, \quad M_A(\eta) = 0 \quad \text{for every } A \subseteq [k+1].$$

The quantity $\Delta_{k,d}(q)$ is its degeneracy probability.

This definition records relations among labelled vertices, so repeated vertices can themselves generate degeneracy. It also excludes all field multiples of the intrinsic relation, not only its two signed integer representatives.

Lemma II.93.2 (Zero-side degeneracy). *If $h_i = 0$ for some i , then the arrangement is degenerate.*

Proof. Choose two labels differing only in ε_i , with the same j . Their vertices coincide. The coefficient array with entries 1, -1 on those labels and 0 elsewhere annihilates all moments. In characteristic two the two nonzero entries are both 1 and still give zero. Its support has size two, whereas every nonzero multiple of e has full support $m \geq 8$. $\square$

II.93.3 Quotienting coefficient arrays before using a union bound

Suppose $\eta' = a\eta + be$ with $a \neq 0$. By (II.93.4), the two arrays have precisely the same simultaneous moment-zero event. We therefore identify them whenever both arrays are admissible. On dividing pointwise by e , this is simply affine equivalence of profiles:

$$t_{\varepsilon,j} = \eta_{\varepsilon,j}/e_{\varepsilon,j}, \quad t' = at + b\mathbf{1}. \quad (\text{II.93.5})$$

Constant profiles correspond exactly to multiples of the intrinsic relation.

If t, t' are nonconstant prime-field profiles and (II.93.5) holds with $a, b \in \mathbb{F}_q$, then in fact $a, b \in \mathbb{F}_p$. Indeed, choose two coordinates where t differs; their differences determine a , and either coordinate then determines b . Thus coefficient-class counts depend on the characteristic p , not on the extension degree of q .

Only arrays with $M_\emptyset(\eta) = 0$ can occur in the exceptional set. The intrinsic signs have equally many positive and negative entries as integer signs. Hence the counting problem is the following one: count nonconstant ternary profiles on an even number of positions with a balanced signed sum, modulo affine equivalence. We solve that problem in Section II.95.

II.94 Boolean coefficient extraction and a rank lemma

Source 27, delivered Section 3.

The original Walsh argument is naturally adapted to odd characteristic. For the present purpose the Boolean zeta transform gives a simpler identity with no division by two.

For $B \subseteq [k]$, define

$$c_{B,j} = \sum_{\varepsilon \in \{0,1\}^k} \eta_{\varepsilon,j} \prod_{i \in B} \varepsilon_i. \quad (\text{II.94.1})$$

For $A \subseteq [k]$, write

$$F_{A,j}(x_j, h) = \sum_{\varepsilon} \eta_{\varepsilon,j} \prod_{i \in A} (x_{ji} + \varepsilon_i h_i), \quad h_B = \prod_{i \in B} h_i.$$

Direct expansion gives

$$F_{A,j} = \sum_{B \subseteq A} h_B c_{B,j} \prod_{i \in A \setminus B} x_{ji}. \quad (\text{II.94.2})$$

The moments in (II.93.3) can be written

$$M_A(\eta) = \sum_j F_{A,j}, \quad M_{A \cup \{k+1\}}(\eta) = \sum_j u_j F_{A,j}. \quad (\text{II.94.3})$$

It is important that *both* equations are available for every A .

Lemma II.94.1 (Parity detection). *For a fixed column j , all $c_{B,j}$ with $B \subsetneq [k]$ vanish if and only if $\eta_{\varepsilon,j}$ is a scalar multiple of $\pi(\varepsilon)$.*

Proof. If the support of a Boolean vertex is $E \subseteq [k]$, then $c_{B,j} = \sum_{E \supseteq B} \eta_{E,j}$. Boolean Möbius inversion gives

$$\eta_{E,j} = \sum_{B \supseteq E} (-1)^{|B \setminus E|} c_{B,j}.$$

If only $c_{[k],j}$ can be nonzero, the right side is $(-1)^{k-|E|} c_{[k],j}$, a scalar parity column. Conversely, a parity column has zero sum over every proper upper interval. The inversion matrix has integer diagonal entries 1, so the argument is valid over every field. $\square$

Call an array *column-parity* if each of its columns is a scalar parity column. An array that is not column-parity has a smallest integer $r < k$ for which $c_A = (c_{A,1}, \dots, c_{A,n}) \neq 0$ for some $|A| = r$. Fix such an A and set $c = c_A$. Equation (II.94.2) gives

$$F_{A,j} = h_A c_j, \quad (\text{II.94.4})$$

and, for every $i \notin A$,

$$F_{A \cup \{i\},j} = h_A c_j x_{ji} + b_{i,j}. \quad (\text{II.94.5})$$

Here $b_{i,j}$ can depend on h and on the coordinates of x_j indexed by A , but not on *any* of the variables $x_{1i}, \dots, x_{ni}$. It can involve other degree- r coefficients, so it must not be replaced by just $h_{A \cup \{i\}} c_{A \cup \{i\},j}$.

Lemma II.94.2 (Two necessary equations). *Assume all $h_i \neq 0$. For a non-column-parity array and the above choice of c , simultaneous vanishing of the moments implies*

$$\sum_j c_j = 0, \quad c \cdot u = 0, \quad (\text{II.94.6})$$

and two affine equations in $z = (x_{1i}, \dots, x_{ni})$ with coefficient rows c and $u \odot c = (u_j c_j)_j$, for each $i \notin A$.

If the u_j are pairwise distinct and $\sum_j c_j = 0$, these two rows have rank two.

Proof. Apply (II.94.3) first to (II.94.4), and then to (II.94.5); divide by the nonzero h_A . A nonzero c of coordinate sum zero has support of size at least two. If the rows were dependent, u_j would have the same value on every index in that support, contrary to pairwise distinctness. $\square$

Proposition II.94.3 (An all-characteristic per-direction bound). *For a fixed non-column-parity coefficient direction, conditional on all $h_i \neq 0$, its moment-zero probability is at most*

$$q^{-2} + (q-1)q^{-n}. \quad (\text{II.94.7})$$

Proof. Use the first nonzero coefficient vector c above. If $\sum c_j \neq 0$, the probability is zero. If $c \notin \mathbb{F}_q s$, the equation $c \cdot u = 0$ cuts a codimension-one subspace of U , costing a factor q^{-1} . The first affine base equation has a nonzero coefficient row and costs another q^{-1} .

If $c = \lambda s \neq 0$, the cross-section equation is automatic. The two base rows have rank two unless all the u_j coincide. The diagonal line lies in U and has relative size $q^{-(n-2)}$. Off the diagonal the probability is at most q^{-2} ; on it the first equation alone gives at most q^{-1} . Averaging gives

$$(1 - q^{-(n-2)})q^{-2} + q^{-(n-2)}q^{-1} = q^{-2} + (q-1)q^{-n}.$$

The other base variables may be fixed throughout, since the two equations are affine in the remaining coordinate block. $\square$

II.95 Counting coefficient directions exactly

Source 27, delivered Section 4.

II.95.1 Balanced profiles

For even $\ell \geq 2$, let $Q_\ell(p)$ count affine classes of nonconstant profiles $t \in D_p^\ell$ satisfying

$$\sum_{i=1}^{\ell/2} t_i - \sum_{i=\ell/2+1}^{\ell} t_i = 0 \quad \text{in } \mathbb{F}_p. \quad (\text{II.95.1})$$

Let $U_\ell(p)$ count affine classes of nonconstant profiles without the balance condition. Define

$$T_\ell(p) = \sum_{a \equiv 0 \pmod{p}} [z^a](z^{-1} + 1 + z)^\ell, \quad (\text{II.95.2})$$

$$B_\ell(p) = \sum_{j \in \mathbb{Z}} \binom{\ell}{\ell/2 + jp}, \quad (\text{II.95.3})$$

where binomial coefficients outside their ordinary range are zero.

Theorem II.95.1 (Affine ternary profile counts). *For $p \geq 5$,*

$$Q_\ell(p) = \frac{T_\ell(p) - 2B_\ell(p) + 1}{2}, \quad (\text{II.95.4})$$

$$U_\ell(p) = \frac{3^\ell - 2^{\ell+1} + 1}{2}. \quad (\text{II.95.5})$$

For $p = 2$ or 3 ,

$$Q_\ell(p) = \frac{p^{\ell-2} - 1}{p - 1}, \quad (\text{II.95.6})$$

$$U_\ell(p) = \frac{p^{\ell-1} - 1}{p - 1}. \quad (\text{II.95.7})$$

If $p > \ell$, then $T_\ell(p) = T_\ell$ and $B_\ell(p) = \binom{\ell}{\ell/2}$.

Proof. First take $p \geq 5$. Changing the signs of the last half of the coordinates is a bijection of the ternary alphabet, so the number of balanced profiles is $T_\ell(p)$. The number of balanced binary profiles is $B_\ell(p)$: if a, b positions are selected in the two halves, balance says $a - b \equiv 0 \pmod{p}$, and Vandermonde's identity sums their number to (II.95.3).

There are three constant ternary profiles. For each of the three unordered pairs of distinct alphabet values, the number of profiles using exactly that pair and satisfying balance is $B_\ell(p) - 2$. All two-valued profiles are affinely equivalent to binary ones, and the only ambiguity of the underlying nontrivial subset is complementation. They therefore give $(B_\ell(p) - 2)/2$ classes.

The number of profiles using all three values is $T_\ell(p) - 3 - 3(B_\ell(p) - 2) = T_\ell(p) - 3B_\ell(p) + 3$. An affine bijection preserving $\{-1, 0, 1\}$ in characteristic at least five is either the identity or negation. To see this, the image of 0 must be the midpoint of the images of -1 and 1 ; the only alphabet element that is the midpoint of the other two is 0 when $p \neq 3$. Thus these profiles form classes of size two. Adding the two types gives (II.95.4). Removing balance replaces $T_\ell(p)$ by 3^ℓ and $B_\ell(p)$ by 2^ℓ , proving (II.95.5).

For $p = 2, 3$, the alphabet is the entire prime field. Balanced profiles form an $(\ell - 1)$ -dimensional vector space containing the constant line. Quotienting by that line gives dimension $\ell - 2$; quotienting nonzero vectors by nonzero scalar multiplication gives (II.95.6). Without balance the quotient dimension is $\ell - 1$, giving (II.95.7). The last assertion follows because the relevant integer sums have absolute value at most ℓ (and the binary differences at most $\ell/2$). $\square$

II.95.2 The two counts needed for arrangements

Set

$$L_{d,p} = Q_n(p), \quad R_{k,d,p} = Q_m(p) - U_n(p). \quad (\text{II.95.8})$$

The first quantity will count cross-section hyperplanes. The second is the number of non-column-parity coefficient directions satisfying the constant moment condition.

Indeed a column-parity array has the form

$$\eta_{\varepsilon,j} = \pi(\varepsilon)c_j, \quad c_j \in D_p. \quad (\text{II.95.9})$$

Its empty moment vanishes automatically, since $k \geq 1$. Dividing by e gives an arbitrary profile c_j/s_j on the n columns, repeated along each column. The resulting number of nonconstant affine classes is $U_n(p)$. All balanced arrays on m positions give $Q_m(p)$ classes; subtracting proves the interpretation of $R_{k,d,p}$. In particular it is a nonnegative integer, even though many of these directions may never vanish on a particular regular region.

For $p > m$, formulas (II.95.8) reduce exactly to (F.32.4) and (F.32.5). For the two small characteristics they become

$$L_{d,p} = \frac{p^{n-2} - 1}{p - 1}, \quad R_{k,d,p} = \frac{p^{m-2} - p^{n-1}}{p - 1}. \quad (\text{II.95.10})$$

Table II.95.1: Some exact large-characteristic constants. Here $p > n2^k$.

d	k	L_d	$S_{k,d}$	$R_{k,d}$
2	1	4	15	459
2	2	4	32	2 585 419
3	1	51	63	35 670
3	2	51	128	13 971 463 734
4	1	484	255	2 582 419
8	1	2 585 444	65 535	79 556 623 497 667

[write] **6 October 2026; sources 21–27.** The entry $R_{2,2} = 2,585,419$ is not $H_8 = C_8 = 2,585,444$ of sources 18 and 25 (they differ by 25, checked at the write); both are correct.

II.96 All first-order obstructions

Source 27, delivered Section 5.

II.96.1 The cross-section hyperplanes

For a column-parity array (II.95.9), all its moments vanish except possibly those indexed by $[k]$ and $[k + 1]$. If all sides are nonzero, these last two moments vanish exactly when

$$\sum_j c_j = 0, \quad c \cdot u = 0. \quad (\text{II.96.1})$$

The excluded intrinsic arrays are precisely $c \in \mathbb{F}_q s$.

Definition II.96.1. Let $\mathcal{H}_{d,p}$ be the collection of distinct hyperplanes of U of the form

$$H_c = \{u \in U : c \cdot u = 0\}, \quad c \in D_p^n, \quad \sum_j c_j = 0, \quad c \notin \mathbb{F}_q s.$$

Define the regular cross-section set and its density by

$$U^\circ = U \setminus \bigcup_{H \in \mathcal{H}_{d,p}} H, \quad g_{d,p}(q) = \frac{|U^\circ|}{q^{n-1}}. \quad (\text{II.96.2})$$

Lemma II.96.2. *The collection $\mathcal{H}_{d,p}$ has exactly $L_{d,p}$ members. Every $u \in U^\circ$ has pairwise distinct coordinates.*

Proof. Two nonzero linear forms on U have the same kernel if and only if they are scalar multiples as forms on U . The annihilator of U is the line $\mathbb{F}_q s$, so $H_c = H_{c'}$ exactly when $c' = ac + bs$ for some $a \neq 0$. Dividing by s turns this into affine profile equivalence. The balance condition becomes a balanced signed sum, so the class count is $Q_n(p) = L_{d,p}$.

If $u_i = u_j$ for distinct indices, choose c supported on those two coordinates with entries $1, -1$. It is balanced and not proportional to s , because $n \geq 4$ and s has full support. Thus $u \in H_c$, contradicting regularity. $\square$

The diagonal translation direction $\mathbf{1}$ is contained in every member of $\mathcal{H}_{d,p}$, because $\sum c_j = 0$. Thus this is essentially a hyperplane arrangement in $U/\mathbb{F}_q \mathbf{1}$, of dimension $n - 2$. No general theory of hyperplane arrangements is required for the bounds below; the intersection structure becomes relevant when higher coefficients are sought.

II.96.2 The sharp first-order theorem in every characteristic

Theorem II.96.3. *For $k \geq 1, d \geq 2$, over any finite field of characteristic p , put $A_p = k + L_{d,p}$. Then*

$$\frac{A_p}{q} - \frac{\binom{A_p}{2}}{q^2} \leq \Delta_{k,d}(q) \leq \frac{A_p}{q} + R_{k,d,p}(q^{-2} + (q-1)q^{-n}). \quad (\text{II.96.3})$$

In particular, for fixed k, d, p and $q = p^a \rightarrow \infty$,

$$\Delta_{k,d}(q) = \frac{k + L_{d,p}}{q} + O_{k,d,p}(q^{-2}). \quad (\text{II.96.4})$$

Proof. There are k side hyperplanes $h_i = 0$ and $L_{d,p}$ cross-section hyperplanes. Every tuple on any of them is degenerate: use Lemma II.93.2 or its defining column-parity relation. These hyperplanes are distinct in the full parameter space. Every pair intersects in codimension two. The first two terms of inclusion–exclusion therefore give the lower bound.

For the upper bound, their union has measure at most A_p/q . Outside that union, a degeneracy witness cannot be column-parity. There are $R_{k,d,p}$ remaining coefficient directions, and Proposition II.94.3 bounds the moment-zero probability of each, conditional on nonzero sides. Dropping the factor $(1 - q^{-1})^k \leq 1$ gives the upper bound. Since $n \geq 4$, its remaining term is $O(q^{-2})$. $\square$

[write] **6 October 2026; sources 21–27.** *Credit and a cross-answer.* For $p > m$ this is source 18’s theorem with its lower bound. In small characteristic it gives the first-order count that source 25 asks for (its Question 4; item II.38): $L_{d,p}$ by Theorem II.95.1, for instance $L_{3,5} = 57$, as in source 25’s Example II.89.2. The classification of the additional components in characteristics two and three is not given.

This proves in a concrete counting sense that all order- q^{-1} obstructions have been identified: the complement of their union contains only $O(q^{D-2})$ degenerate tuples. The word “first-order” here refers to this counting statement; no unproved assertion about geometric irreducibility is being used.

Example II.96.4 (Characteristic dependence). For $d = 3$, the numbers $L_{3,p}$ are

p	2	3	5	$p > 6$
$L_{3,p}$	15	40	57	51.

In characteristic five, ternary integer sums equal to ± 5 also become zero, increasing the count. Large field size alone does not remove this effect: it persists in every extension of $\mathbb{F}_5$.

II.97 Second-order structure on the regular region

Source 27, delivered Section 6.

From now until the end of Section II.98, assume $p > m$. Define

$$G = (1 - q^{-1})^k g_{d,p}(q). \quad (\text{II.97.1})$$

This is exactly the probability that all sides are nonzero and $u \in U^\circ$. Its complement is entirely degenerate.

II.97.1 Directions that are not second-order

Fix a non-column-parity direction and let $r < k$ be its smallest nonzero Boolean coefficient degree, as in Section II.94. If some degree- r vector c_A is not proportional to s , then its event on the regular region costs one cross-section equation and two base equations. More precisely, $c_A \cdot u = 0$ defines a codimension-one subspace in U ; after fixing such a regular u , Lemma II.94.2 gives rank two in any coordinate block outside A . Hence its probability on the regular region, conditional only on nonzero sides, is at most q^{-3} .

Suppose instead that every nonzero degree- r vector is proportional to s . If $r \leq k - 2$, choose one such A and two different indices $i, \ell \notin A$. The paired moments for $A \cup \{i\}$ give two affine equations in the i th base-coordinate block; those for $A \cup \{\ell\}$ give two in the ℓ th block. Their right sides depend only on the A -coordinates and h , not on either other variable block. On regular u , each pair has rank two. Together they have rank four and cost q^{-4} .

Thus the only remaining possibility for a direction contributing at the second order is

$$c_{B,j} = 0 \ (|B| \leq k - 2), \quad c_{[k] \setminus \{i\},j} = \lambda_i s_j \ (1 \leq i \leq k), \quad (\text{II.97.2})$$

with at least one $\lambda_i \neq 0$.

Definition II.97.1. A non-column-parity direction satisfying (II.97.2) is called a *core direction*. This term is local to this article.

Proposition II.97.2. *For every non-column-parity direction that is not a core direction, its moment-zero event on the regular region, conditional on nonzero sides, has probability at most q^{-3} .*

Proof. The two cases above exhaust the possibilities. The first gives q^{-3} . The second gives at most $g_{d,p}(q)q^{-4} \leq q^{-3}$. $\square$

II.97.2 Classification of all core profiles

Lemma II.97.3 (Affine Boolean form). *Condition (II.97.2) is equivalent to*

$$\eta_{\varepsilon,j} = e_{\varepsilon,j} \left(t_j + \sum_{i=1}^k b_i \varepsilon_i \right) \quad (\text{II.97.3})$$

for some $t_j, b_i \in \mathbb{F}_p$, with $b \neq 0$.

Proof. Divide each column by its parity and by s_j . The Boolean coefficient transform is invertible. Vanishing of the coefficients below degree $k - 1$ means that the resulting Boolean function has degree at most one, and proportionality of its degree- $k - 1$ coefficients to s_j says that the linear coefficients are common to the columns.

For completeness this can also be checked directly: a constant term times parity has no proper nonzero c_B ; a term ε_i times parity has no nonzero c_B unless $B \supseteq [k] \setminus \{i\}$, and at that set its coefficient is $(-1)^{k-1}$. These $k + 1$ arrays are linearly independent. Boolean inversion shows that they exhaust the solution space of the stated vanishing conditions in each column. The common proportionality factors then give common b_i . $\square$

Every value in parentheses in (II.97.3) must belong to the three-element alphabet. This severely restricts b .

Lemma II.97.4 (Width two). *If $p > m$ and a nonconstant affine Boolean function takes values in $\{-1, 0, 1\} \subset \mathbb{F}_p$, its slopes have integer representatives b_i with*

$$\sum_i |b_i| \leq 2. \quad (\text{II.97.4})$$

Consequently it has either one active coordinate, or two active coordinates with unit slopes.

Proof. Choose the representative $t \in \{-1, 0, 1\}$ of the value at the zero vertex. Each slope is a difference of two alphabet values, so has a representative in $\{-2, -1, 0, 1, 2\}$. At any vertex the integer affine sum differs from the chosen representative of its field value by an integer of absolute value at most $2k + 2$. Since $p > m \geq 4 \cdot 2^k > 2k + 2$, no nonzero multiple of p can occur. Thus the ordinary integer affine function really takes all its values in $[-1, 1]$. Its range width is $\sum_i |b_i|$, proving the claim. $\square$

Remark II.97.5 (The characteristic cutoff in the width argument). The proof actually requires only $p > 2k + 2$ (and hence $p \geq 5$), not $p > m$. Accordingly the core count and the resolved bound remain valid in that range with $L_{d,p}, R_{k,d,p}$ in place of their clean large-characteristic forms. The stronger assumption $p > m$ in Theorem F.32.1 also removes modular wraparound from the balanced coefficient-class counts.

Theorem II.97.6 (The core direction count). *There are exactly*

$$S = S_{k,d} = k(2^n - 1) + 2 \binom{k}{2}$$

core directions. For fixed nonzero sides and regular u , every one of their moment-zero events has probability exactly q^{-2} in the base variables.

Proof. For one active coordinate i , divide by its slope and use translation of the entire profile. A unit slope allows each column offset to be one of two adjacent values. A nonconstant binary offset pattern gives one direction; there are $2^n - 2$ such patterns. The two constant patterns are equivalent under translation and contribute one further direction. A slope of absolute value two forces a constant offset and belongs to that same constant-pattern direction after affine normalization. This gives $2^n - 1$ directions for each i .

For two active coordinates i, ℓ , their slopes have absolute value one. After common scaling, their ratio is either 1 or -1 . The total width is two, so every column offset is forced to the same value: -1 for slopes $(1, 1)$, or 0 for slopes $(1, -1)$. These give two directions for every unordered coordinate pair. Distinct slope ratios, coordinate supports, and nonconstant offset patterns cannot be related by one common affine transformation. This proves the count.

For a core profile, all moments except the two top ones vanish identically on U . Expanding the top moment shows that the remaining conditions are

$$\sum_j s_j \left(t_j + \sum_i b_i \frac{x_{ji}}{h_i} \right) = 0, \quad \sum_j s_j u_j \left(t_j + \sum_i b_i \frac{x_{ji}}{h_i} \right) = 0. \quad (\text{II.97.5})$$

To check the expansion, each ε_i term contributes the factor $1 + x_{ji}/h_i$ after division by the common alternating top-cube factor. The additional constant $\sum_i b_i$ disappears from both sums because $\sum s_j = \sum s_j u_j = 0$.

At least one b_i is nonzero. In that coordinate block the coefficient rows are nonzero multiples of s and $s \odot u$. Since the u_j are pairwise distinct, these rows are independent. Therefore (II.97.5) is an affine rank-two system and has exactly a q^{-2} fraction of solutions. $\square$

[write] **7 October 2026; sources 40–46.** Source 44’s Lemma II.175.1 removes the condition $p > 2k + 2$: the core directions are those of this theorem in every characteristic $p \geq 5$, and the union of the core events is counted exactly (its Theorems II.176.2 and II.176.3).

II.97.3 Intersections of core events

Lemma II.97.7. *Fix regular u and nonzero sides. Distinct core directions with parallel slope vectors have disjoint zero events. Core directions with nonparallel slope vectors have intersection probability q^{-4} .*

Proof. Only the one-coordinate families have more than one direction with a given slope. Normalize both to slope 1. If their zero events meet, subtracting (II.97.5) gives

$$\sum_j s_j (t_j - t'_j) = 0, \quad \sum_j s_j u_j (t_j - t'_j) = 0.$$

The differences can be chosen in $\{-1, 0, 1\}$. For distinct directions the vector $c_j = s_j(t_j - t'_j)$ is not proportional to s , since a constant difference would make the profiles affinely equivalent. Thus u belongs to a forbidden hyperplane, a contradiction.

For nonparallel slopes b, b' , consider the two independent coordinate functionals b_i/h_i and b'_i/h_i , and the two independent column functionals s and $s \odot u$. The four tensor-product functionals are linearly independent. They are exactly the coefficient rows of the four affine equations defining the intersection, so that system is surjective onto $\mathbb{F}_q^4$ and has fraction q^{-4} of solutions. $\square$

Put

$$P_{k,d} = \binom{S_{k,d}}{2} - k \binom{2^n - 1}{2}. \quad (\text{II.97.6})$$

This is the number of unordered pairs of nonparallel core directions. Bonferroni's inequality and Lemma II.97.7 show that the conditional probability of their union lies between

$$\frac{S_{k,d}}{q^2} - \frac{P_{k,d}}{q^4} \quad \text{and} \quad \frac{S_{k,d}}{q^2}. \quad (\text{II.97.7})$$

There is no q^{-3} loss from intersections of core events on this region.

Theorem II.97.8 (Hyperplane-resolved bound). *Let $p > m$ and use G from (II.97.1). Then*

$$\Delta_{k,d}(q) \geq 1 - G + G \left(\frac{S_{k,d}}{q^2} - \frac{P_{k,d}}{q^4} \right), \quad (\text{II.97.8})$$

$$\Delta_{k,d}(q) \leq 1 - G + G \frac{S_{k,d}}{q^2} + (1 - q^{-1})^k \frac{R_{k,d} - S_{k,d}}{q^3}. \quad (\text{II.97.9})$$

Proof. The complement of the regular region has measure $1 - G$ and is entirely degenerate. Apply (II.97.7) on the regular region. For the upper bound add the $R_{k,d} - S_{k,d}$ remaining directions, each bounded by Proposition II.97.2. Their estimates were conditional on nonzero sides, so the factor $(1 - q^{-1})^k$ is retained. $\square$

Proof of Theorem F.32.1. The lower bound is the lower bound of Theorem II.96.3. For the upper bound in (II.97.9), use $1 - G \leq (k + L_d)/q$, $G \leq 1$, and $(1 - q^{-1})^k \leq 1$. This proves (F.32.7). Its lower and upper bounds have the same q^{-1} coefficient, proving sharpness. $\square$

II.98 Two pairs of cross sections and the second coefficient

Source 27, delivered Section 7.

For $d = 2$, $s = (1, 1, -1, -1)$. Subtract a common translation from u and write

$$u = (a, b, 0, a + b). \quad (\text{II.98.1})$$

The quotient $U/\mathbb{F}_q \mathbf{1}$ is thus identified with the (a, b) -plane.

Proposition II.98.1. *In odd characteristic,*

$$g_{2,p}(q) = \frac{(q-1)(q-3)}{q^2} = 1 - \frac{4}{q} + \frac{3}{q^2}. \quad (\text{II.98.2})$$

In characteristic two,

$$g_{2,2}(q) = \frac{(q-1)(q-2)}{q^2}. \quad (\text{II.98.3})$$

Proof. For odd characteristic, $L_{2,p} = 4$ by Theorem II.95.1. The four hyperplanes are represented in the quotient by the distinct lines

$$a = 0, \quad b = 0, \quad a = b, \quad a = -b.$$

They can all be obtained from pairwise cross-section equalities. Their union contains $1 + 4(q-1)$ points. Its complement therefore has $q^2 - 4q + 3 = (q-1)(q-3)$ points. In characteristic two the last two lines coincide, and $L_{2,2} = 3$; the same argument gives $q^2 - 3q + 2 = (q-1)(q-2)$. $\square$

Proof of the asymptotic assertion in Theorem F.32.3. Under its large-characteristic assumption,

$$G = (1 - q^{-1})^k(1 - 4q^{-1} + 3q^{-2}) = 1 - \frac{k+4}{q} + \frac{\binom{k}{2} + 4k + 3}{q^2} + O_k(q^{-3}).$$

Here $S_{k,2} = k^2 + 14k$. Both sides of Theorem II.97.8 therefore have second coefficient

$$S_{k,2} - \binom{k}{2} - 4k - 3 = \frac{k^2 + 21k - 6}{2}.$$

The residual upper term is $O_k(q^{-3})$ and the core-intersection lower term is $O_k(q^{-4})$, proving (F.32.10). $\square$

For $k = 1$, all 15 core directions have parallel slopes, so their events are pairwise disjoint on the regular region. Consequently

$$1 - G + \frac{15G}{q^2} = \frac{5}{q} + \frac{8}{q^2} - \frac{72}{q^3} + \frac{105}{q^4} - \frac{45}{q^5} \quad (\text{II.98.4})$$

is an exact lower bound. The difference from the full degeneracy fraction is between 0 and $444(1 - q^{-1})q^{-3}$. In particular the conventional proofs already give

$$\frac{5}{q} + \frac{8}{q^2} - \frac{72}{q^3} \leq \Delta_{1,2}(q) \leq \frac{5}{q} + \frac{8}{q^2} + \frac{372}{q^3}. \quad (\text{II.98.5})$$

The exact computation below determines the entire remaining expression.

II.99 Two extensions of the second-order theorem

Source 27, delivered Section 8.

II.99.1 A hyperplane statistic determines the second coefficient

Let $\mathcal{L}_2$ be the set of distinct codimension-two subspaces of U that arise as intersections of two members of $\mathcal{H}_{d,p}$. For $F \in \mathcal{L}_2$, let ν_F be the number of members containing F , and put

$$b_{2,d,p} = \sum_{F \in \mathcal{L}_2} (\nu_F - 1). \quad (\text{II.99.1})$$

This is a finite invariant of the cross-section hyperplane arrangement over $\mathbb{F}_p$.

Lemma II.99.1. *For fixed d, p and $q = p^a$,*

$$g_{d,p}(q) = 1 - \frac{L_{d,p}}{q} + \frac{b_{2,d,p}}{q^2} + O_{d,p}(q^{-3}). \quad (\text{II.99.2})$$

Proof. Apply finite inclusion–exclusion to the hyperplanes. A nonempty intersection of codimension r has relative size q^{-r} . A codimension-two flat F contributes the alternating sum over subsets of at least two of the ν_F hyperplanes containing it. Any two of these hyperplanes already intersect in F . The resulting coefficient is

$$\sum_{j=2}^{\nu_F} (-1)^j \binom{\nu_F}{j} = \nu_F - 1.$$

All other terms have codimension at least three. Their number is finite and depends only on the original prime-field arrangement. This proves the expansion, and also shows that the exact complement density is a polynomial in q^{-1} for fixed d, p . $\square$

Corollary II.99.2 (The second coefficient for arbitrary d). *Fix k, d and a prime $p > m$. Along the extensions $q = p^a$,*

$$\Delta_{k,d}(q) = \frac{k + L_d}{q} + \frac{S_{k,d} - \binom{k}{2} - kL_d - b_{2,d,p}}{q^2} + O_{k,d,p}(q^{-3}). \quad (\text{II.99.3})$$

Proof. Substitute (II.99.2) in Theorem II.97.8 and expand $(1 - q^{-1})^k$ to second order. The core union has no intersection loss before order q^{-4} , and the remaining directions cost at most q^{-3} . $\square$

[write] **7 October 2026; sources 40–46.** *The statistic $b_{2,d,p}$ is now computed* (source 44, Section II.173): it is characteristic-independent for $p > \max(2d, 16)$ and given for every d by a nine-point grid formula; $b_{2,8} = 3,280,234,631,541$ (Theorems II.178.1, II.179.1, II.180.1).

This gives the second coefficient as a finite linear-algebra problem even when a closed form for $b_{2,d,p}$ is unavailable. One must not silently replace $b_{2,d,p}$ by a characteristic-independent number: ranks among ternary normals may change in particular characteristics. For $d = 2$ in odd characteristic all four quotient lines meet at the origin, so $b_{2,2,p} = 3$, recovering (F.32.10). Hyperplane arrangements and their intersection posets provide a natural language for further computations; a standard background reference is Stanley’s survey [99]. The elementary argument above is the only input needed here.

II.99.2 An exact core count in characteristics two and three

The width-two argument was needed because a large prime field contains many values outside the allowed alphabet. In characteristic two or three, the alphabet is the whole prime field. In these cases the core count is easier, not harder.

Theorem II.99.3 (Small-characteristic cubic bounds). *For $p \in \{2, 3\}$, define*

$$\begin{aligned} B_{k,p} &= \frac{p^k - 1}{p - 1}, & S_{k,d,p} &= p^{n-1} B_{k,p}, \\ P_{k,d,p} &= \binom{S_{k,d,p}}{2} - B_{k,p} \binom{p^{n-1}}{2}. \end{aligned} \quad (\text{II.99.4})$$

The bounds in Theorem II.97.8 remain valid after replacing $S_{k,d}$ by $S_{k,d,p}$, $P_{k,d}$ by $P_{k,d,p}$, and $R_{k,d}$ by $R_{k,d,p}$. In particular, for every extension $q = p^a$,

$$\Delta_{k,d}(q) \leq \frac{k + L_{d,p}}{q} + \frac{S_{k,d,p}}{q^2} + \frac{R_{k,d,p} - S_{k,d,p}}{q^3}. \quad (\text{II.99.5})$$

The second-coefficient formula in (II.99.3) also holds with these characteristic-dependent constants.

Proof. The Boolean form (II.97.3), the rank argument of Proposition II.97.2, and the core equations (II.97.5) are valid in every characteristic. In the present two cases all choices $t \in \mathbb{F}_p^n$ and $0 \neq b \in \mathbb{F}_p^k$ are

admissible. Choose a projective slope direction, of which there are $B_{k,p}$. After fixing a representative of that direction, the offsets are identified exactly by common translation; their quotient has p^{n-1} elements. This proves the core count.

The difference of two offsets is still alphabet-valued, since the alphabet is $\mathbb{F}_p$. The proof of disjointness for parallel slopes in Lemma II.97.7 therefore applies unchanged. Nonparallel slopes again give four independent tensor-product rows. The same union bounds prove the assertions. $\square$

Example II.99.4. For $d = 2, k = 1$, the second-order asymptotics in the two small characteristics are

$$\begin{aligned}\Delta_{1,2}(2^a) &= \frac{4}{2^a} + \frac{3}{2^{2a}} + O(2^{-3a}), \\ \Delta_{1,2}(3^a) &= \frac{5}{3^a} + \frac{20}{3^{2a}} + O(3^{-3a}).\end{aligned}$$

Indeed the hyperplane coefficients are $(L, b_2) = (3, 2)$ and $(4, 3)$, respectively, and the core counts are 8 and 27. These differ from the large-characteristic coefficient 8 in (F.32.11); keeping track of characteristic is mathematically necessary.

II.100 An exact count for eight labelled vertices

Source 27, delivered Section 9.

This section proves the exact assertion in Theorem F.32.3. Its only computer-assisted component is a finite arrangement of rational affine lines, specified below and in the accompanying certificate. The general theorems of the preceding sections do not depend on it.

II.100.1 Normalization on the regular region

Take $k = 1, d = 2$ and characteristic $p \geq 11$. On the regular region, $h \neq 0$ and $u_2 - u_3 \neq 0$. Independent affine changes in the first and final coordinates preserve the span of the squarefree monomials and hence preserve degeneracy. We may normalize

$$h = 1, \quad x_4 = 0, \quad u = (r, 1, 0, r + 1), \quad r \notin \{0, 1, -1\}. \quad (\text{II.100.1})$$

The remaining free variables are $x_1, x_2, x_3 \in \mathbb{F}_q$.

For each fixed normalized tuple there are exactly

$$(q - 1) \cdot q \cdot (q - 1) \cdot q = q^2(q - 1)^2 \quad (\text{II.100.2})$$

original tuples: choose the nonzero side, the first-coordinate translation, the nonzero cross-section scale, and the cross-section translation. This count is independent of the normalized tuple.

In the order $(j, \varepsilon) = (1, 0), (1, 1), \dots, (4, 1)$, put

$$c_j = \eta_{0,j} + \eta_{1,j}, \quad b_j = \eta_{1,j}. \quad (\text{II.100.3})$$

There are exactly 484 balanced non-intrinsic coefficient classes, of which 25 are column-parity and 459 are not. Of the latter, 15 are core. The first two moment constraints for a non-parity array are

$$\sum_j c_j = 0, \quad (c_1 + c_4)r + (c_2 + c_4) = 0.$$

The two remaining constraints are precisely

$$c_1x_1 + c_2x_2 + c_3x_3 = -\sum_j b_j, \quad (\text{II.100.4})$$

$$c_1rx_1 + c_2x_2 = -(b_1 + b_4)r - b_2 - b_4. \quad (\text{II.100.5})$$

For a regular ratio and an eligible non-parity class these two equations have rank two, so they describe an affine line in $\mathbb{F}_q^3$.

A core class has $c \in \mathbb{F}_q s$ and is eligible for every regular r . For any other class, at most one ratio can be eligible, unless the cross-section equation is a nonzero constant and hence impossible. The complete list of rational exceptional ratios is

$$\mathcal{E} = \{-3, -2, -\frac{1}{2}, -\frac{1}{3}, \frac{1}{3}, \frac{1}{2}, 2, 3\}. \quad (\text{II.100.6})$$

Outside this list there are only the 15 disjoint core lines, containing $15q$ normalized degenerate base triples.

II.100.2 The finite line certificate

To make the finite calculation unambiguous, the complete algorithm is as follows. All operations in this paragraph take place over $\mathbb{Q}$.

1. Enumerate the 3^8 vectors $\eta \in \{-1, 0, 1\}^8$ and retain those with integer coordinate sum zero. Divide coordinatewise by $e = (1, -1, 1, -1, -1, 1, -1, 1)$ to obtain t .
2. Discard constant t . Canonicalize every other profile by subtracting its first entry and dividing by its first nonzero resulting entry. Keep one η for every canonical profile. Exactly 484 profiles remain.
3. Discard the 25 column-parity profiles. For each of the remaining 459, calculate c, b from (II.100.3). Those with $c_1 + c_4 = c_2 + c_4 = 0$ are the 15 core profiles. For every other one, solve $(c_1 + c_4)r + c_2 + c_4 = 0$ and discard impossible or nonregular ratios.
4. For every ratio in (II.100.6), adjoin all 15 core profiles to its eligible profiles. Put the two augmented rows of (II.100.4) and (II.100.5) in reduced row-echelon form. Identical row spaces represent the same affine line and are deduplicated.
5. For every pair of distinct lines, stack their augmented equations. A consistent system of coefficient rank three yields one intersection point. Deduplicate these points and record *all* lines through each point. If μ_P lines meet at P , the correction at that point is $\mu_P - 1$.

The exact output is given in Table II.100.1. The JSON certificate records every rational line equation, every intersection point, and its incident line indices; it is not merely a list of the aggregate numbers.

Lemma II.100.1 (Stability of the certificate). *For every prime $p \geq 11$, reduction of the rational data in Table II.100.1 modulo p preserves the distinct ratios, the distinct lines, every pairwise intersection type, and all intersection-point multiplicities. The same is true over every extension of $\mathbb{F}_p$.*

Table II.100.1: The complete rational line enumeration for the exceptional ratios. The last column is $I_r = \sum_P(\mu_P - 1)$.

r	Additional eligible profiles	Distinct lines N_r	Intersection points	I_r
-3	4	19	18	18
-2	24	39	32	44
-1/2	24	39	32	44
-1/3	4	19	18	18
1/3	4	19	18	18
1/2	24	39	32	44
2	24	39	32	44
3	4	19	18	18
Totals used in the proof:		$\sum_r(N_r - 15) = 112$	$\sum_r I_r = 248$	

Proof by the supplied finite certificate. The code reconstructs the full data directly from the 3^8 coefficient vectors. For every pair of lines it checks both the coefficient matrix and the augmented matrix. In each matrix it chooses a nonzero maximal-rank minor. This amounts to 7296 exact matrix checks across the eight ratios. It also checks all denominators of line and point coordinates, and for every pair of distinct recorded intersection points chooses one nonzero coordinate difference. Finally it checks the differences among all exceptional ratios and $0, 1, -1$.

The prime factors of the selected nonzero minors, nonzero differences, and denominators are all contained in $\{2, 3, 5, 7\}$. This is verified by exact integer factorization in the supplied standard-library program. Therefore none vanishes in characteristic at least 11. All larger minors are zero over $\mathbb{Q}$ and remain zero after reduction; the selected minors keep the rank from dropping. Coefficient and augmented ranks together preserve whether two lines are parallel, skew, or intersecting. The point-difference checks keep distinct intersection points from merging.

No new point-line incidence can then appear: if a point from two lines lay on a third line after reduction, that third line would have to meet one of the first two. Rank preservation forces a rational intersection already to exist, and the distinct-point checks force it to be the original common point. This proves preservation of the complete incidence data. Ranks and distinctness are unchanged by field extension.

There is also no loss of coefficient profiles in reducing this certificate. For $p > 8$ the balanced class count is 484, exactly as over $\mathbb{Q}$. Rational affine equivalences remain valid modulo p (their normalizing denominators divide 2); the equal class counts therefore exclude any new identifications. The coefficients $c_1 + c_4, c_2 + c_4$ have absolute value at most 4, so the classification of their linear equations introduces no new exceptional ratios or discarded regular ratios for $p \geq 11$. The explicit ratio checks exclude collisions in the listed ratios themselves. $\square$

Remark II.100.2 (Scope of the computer-assisted step). The lemma is a finite, reproducible computation, not an appeal to a numerical pattern. Its rational input space has 6561 elements; every elimination, determinant, incidence test, and factorization is exact. The code is included so that the table and its exceptional-prime assertion can be audited or formalized. It has not been checked by a proof assistant.

II.100.3 Counting the union and restoring the parameters

Distinct affine lines in three-dimensional affine space meet in at most one point. If their distinct intersection points are P , the size of their union is exactly

$$N_r q - \sum_P (\mu_P - 1) = N_r q - I_r. \quad (\text{II.100.7})$$

This formula already accounts for multiple intersections; subtracting only the number of intersecting pairs would be incorrect when three or more lines concur.

There are $q-3$ regular ratios. For $p \geq 11$ all eight exceptional ratios are distinct and regular. Start with the $15q$ contribution at each regular ratio, and then add the corrections from Table II.100.1. The total number of normalized degenerate triples, summed over regular ratios, is

$$M(q) = 15q(q-3) + 112q - 248 = 15q^2 + 67q - 248. \quad (\text{II.100.8})$$

Since $G = (1 - q^{-1})(1 - 4q^{-1} + 3q^{-2})$, the full degeneracy proportion is

$$\Delta_{1,2}(q) = 1 - G + \frac{q^2(q-1)^2 M(q)}{q^8} = 1 - G + \frac{(q-1)^2}{q^6} M(q). \quad (\text{II.100.9})$$

Expansion proves (F.32.11). Equivalently, the integer count is

$$5q^7 + 8q^6 + 40q^5 - 367q^4 + 563q^3 - 248q^2. \quad (\text{II.100.10})$$

This completes the proof of the exact part of Theorem F.32.3.

II.100.4 An independent subset-sum test

There is a different way to test degeneracy without using the line classification. For each labelled vertex, form its column of squarefree moments. Let V_S be the sum of the columns indexed by a subset S of the m labels. Every collision $V_S = V_T$ yields a ternary relation by taking the difference of the two indicator vectors.

Lemma II.100.3 (Subset-sum criterion). *In odd characteristic, an arrangement with m labels is nondegenerate if and only if its 2^m subset moment sums have exactly $2^m - 1$ distinct values.*

Proof. The positive and negative supports of e give one unavoidable collision. Conversely every ternary relation is a difference of two subset indicator vectors. If all such relations are intrinsic, their nonzero coefficient vectors can only be e and $-e$: an intrinsic scalar multiple has all coordinates nonzero, and an alphabet-valued one in odd characteristic has scalar 1 or -1 . Full support means that the two subsets giving this collision must be exactly the positive and negative supports of e , in either order. Thus there is exactly one duplicated value and all other subset sums are distinct. Any further collision supplies a non-intrinsic relation. $\square$

For the present eight-vertex case the four moment columns are $(1, x, u, xu)$. The accompanying independent test enumerates every normalized base triple for every regular ratio at $q = 5, 7, 11, 13, 17, 19$, calculates all 256 subset sums, and checks whether their number of distinct values is less than 255. Its results are in Table II.100.2.

These 212 766 exhaustive checks corroborate the formula but do not prove its validity for untested fields. That validity comes from the exact rational incidence calculation and Lemma II.100.1, which also cover all extension fields of the permitted characteristics.

Table II.100.2: Independent exhaustive checks, not using the line certificate. The failed predictions at 5, 7 are intentional negative controls outside the exact theorem’s characteristic range.

q	Base triples examined	Observed $M(q)$	$15q^2 + 67q - 248$	Match
5	250	250	462	No
7	1 372	900	956	No
11	10 648	2 304	2 304	Yes
13	21 970	3 158	3 158	Yes
17	68 782	5 226	5 226	Yes
19	109 744	6 440	6 440	Yes

II.101 Consequences for an exceptional-arrangement budget

Source 27, delivered Section 10.

II.101.1 Absolute and relative counting

The benefit to Gowers’s selection step is a smaller sufficient ambient field size for making the number of exceptional arrangements negligible. The printed application of Lemma 15.4 in the proof of Lemma 15.5 charges its full ambient degeneracy bound against a positive expectation surplus [1, p. 565]. The present estimates can be inserted wherever that same absolute count is the only use of Lemma 15.4.

Proposition II.101.1 (Relative and weighted versions). *Let $\mathcal{A}$ be any collection of arrangement tuples with $|\mathcal{A}| \geq \rho q^D$, where $\rho > 0$. The fraction of degenerate tuples in $\mathcal{A}$ is at most $\Delta_{k,d}(q)/\rho$.*

More generally, let w be a nonnegative weight on ambient tuples with $w \leq W$ and $\sum w \geq \rho q^D$. Then the fraction of total weight on degenerate tuples is at most

$$\frac{W \Delta_{k,d}(q)}{\rho}. \quad (\text{II.101.1})$$

Proof. There are $\Delta_{k,d}(q)q^D$ ambient degenerate tuples. Their contribution in a subset is no larger, and their weighted contribution is at most W times their number. Divide by the stated lower bound on total mass. $\square$

Thus to guarantee an exceptional fraction at most ε inside a collection of relative density at least ρ , it suffices to use the ambient budget $\rho\varepsilon$ in Corollary F.32.2. In the weighted setting use $\rho\varepsilon/W$. This is an absolute-to-relative transfer, not an assertion that the hyperplane events remain equidistributed within an arbitrary structured subset.

II.101.2 Quantitative comparison

The original estimate (F.32.2) requires $q \geq k3^m/\varepsilon$ for an ambient error budget ε . In the new estimate, only the third term contains the full coefficient-direction inventory. Since $R_{k,d} \leq 3^m/2$, that contribution is at most

$$\left(\frac{3^{m+1}}{2\varepsilon} \right)^{1/3}$$

as a sufficient field-size threshold. This statement does not approximate $R_{k,d}$; it uses only a crude upper bound for it. In particular the exponential dependence of that *particular* threshold changes from 3^m to at most a constant times $3^{m/3}$. The separate hyperplane threshold can still dominate.

Table II.101.1: Sufficient real thresholds for $\varepsilon = 0.01$, rounded up in the last column. A field with the required characteristic and size must still be chosen. No assertion of primality is made for the displayed integers.

d	k	Threshold from $k3^m/\varepsilon$	Threshold from (F.32.9)
2	1	656 100	1 500
2	2	8 609 344 200	1 800
3	1	53 144 100	15 600
8	1	185 302 018 885 184 100	775 633 500

For the $d = 8$ arrangements used in the selection argument, the first-order coefficient is

$$k + 2\,585\,444, \quad (\text{II.101.2})$$

rather than $k3^{16 \cdot 2^k}$. The precise cubic remainder and the characteristic requirement must accompany this coefficient in any finite application; quoting the leading term alone would be insufficient.

II.101.3 Limits of the conclusion

A complete improved bound for Szemerédi’s theorem would require rerunning all subsequent parameter choices, preserving every hypothesis, and checking whether another threshold dominates. This article has not done that. Neither Table II.101.1 nor (II.101.2) is a claim about the best known global density bound or about a final improvement to Gowers’s Theorem 18.2. The contribution is an explicit replacement for the local exceptional count and a documented interface for carrying that replacement further.

II.102 Boundary cases and safeguards

Source 27, delivered Section 11.

II.102.1 Why the hypothesis $d \geq 2$ matters

When $d = 1$, U is the diagonal line $u_1 = u_2$. The two rows s and $s \odot u$ are always dependent. Therefore the rank-two gain underlying Proposition II.94.3 is not available.

Proposition II.102.1 (The endpoint $k = d = 1$). *In odd characteristic,*

$$\Delta_{1,1}(q) = \frac{4}{q} - \frac{3}{q^2}. \quad (\text{II.102.1})$$

In characteristic two it is $3/q - 2/q^2$.

Proof. All final coordinates agree. The four first-coordinate labels are $x_1, x_1 + h, x_2, x_2 + h$, and the intrinsic relation is $(1, -1, -1, 1)$. Up to affine profile equivalence, the non-intrinsic balanced ternary relations yield exactly

$$h = 0, \quad x_1 - x_2 = 0, \quad x_1 - x_2 = h, \quad x_1 - x_2 = -h.$$

For an explicit exhaustion argument, Theorem II.95.1 gives $Q_4(p) = 4$ for every odd characteristic. Balanced non-intrinsic witnesses for the four displayed lines belong to four distinct classes, so these exhaust all possibilities. They are four distinct central lines in the $(h, x_1 - x_2)$ -plane and have union size $1 + 4(q - 1)$. The common translation and final coordinate cancel when taking the fraction. In characteristic two the last two lines coincide; the same count uses three lines. $\square$

[write] **6 October 2026; sources 21–27.** *Credit.* In odd characteristic this is source 22’s Theorem II.85.1 at $k = 1$; source 22 proves the leading coefficient $(k^2 + 3k)/p$ for every k , which answers the leading-coefficient part of this source’s question on the single-pair regime (item II.38).

The endpoint has leading coefficient four in odd characteristic, not the one that a naive substitution $L_1 = 0$ into the main theorem would suggest. This is a genuine change of rank behavior, not a removable technicality.

II.102.2 All squarefree moments are required

After Lemma 15.4, Gowers remarks that his proof uses only the full top monomial among the monomials containing the last coordinate [1, p. 563]. The present refinement is deliberately for the *original all-moments definition*, also encoded in the inspected repository interface.

The paired constraints (II.94.3) are indispensable here. For example, when $k = 1$, the relation with coefficient 1 on $(1, 0)$ and -1 on $(2, 0)$ automatically satisfies the constant moment. Its first and mixed moments vanish when $x_1 = x_2 = 0$, even if $u_1 \neq u_2$. The omitted final-coordinate moment would be $u_1 - u_2$, which is nonzero. Thus removing that moment changes the individual witness zero sets and defeats the specific rank argument. No extension to that weakened definition is claimed.

II.102.3 Other distinctions that must be retained

The field size q and the characteristic p play different roles. Counts of prime-field coefficient directions can change with p even when q is arbitrarily large. Modular scalar multiples of the intrinsic relation must be removed before counting; integer signs alone are not enough. Labelled parameter counts must not be divided by cube symmetries without analyzing stabilizers. Finally, the cubic bounds are upper bounds on unions: they do not assert that all non-core directions have distinct or nonempty codimension-three loci.

II.103 Common base partitions and exact transport of affine fibres

Source 23, delivered Section 8.

The interpolation step of Lemma 16.10 has two distinct requirements. Its sampled cross-sections need covers on one common base partition, and each resulting base box must be extended to boxes in the full product. The second requirement is substantive: the final axis of the original product need not have the step of a refined base box. We prove an explicit construction which addresses

both requirements. A preliminary localization from width m to width at least $m/3$ is sufficient; after that localization, transport preserves the complete width of every refined base box. There is no additional power loss at the transport stage. The final axes are allowed to be long, which is appropriate once the fibre functions are already affine on the original final axis.

Throughout this section, p is prime. A proper progression of length n is

$$P(a, d, n) = \{a + jd : 0 \leq j < n\} \subseteq \mathbb{F}_p,$$

with distinct displayed points. A proper k -box is a Cartesian product of k such progressions with one common step; its width is the minimum of the displayed lengths. We assume $k \geq 1$. A multiaffine function on a box means the restriction of a polynomial over $\mathbb{F}_p$ having degree at most one in each variable. This is the convention called multilinearity in the relevant Section 16 interpolation statements.

II.103.1 A sharp elementary localization bound

Lemma II.103.1 (Balanced localization). *Let $m \geq 4$ be an integer, and set*

$$B_m = \lfloor m/2 \rfloor + 1, \quad L_m = \lfloor (2B_m + 1)/3 \rfloor.$$

Every integer interval of length $n \geq m$ has a partition into intervals whose lengths lie in $[L_m, B_m]$. Moreover, $L_m \geq m/3$ and $L_m \geq 2$. Consequently every proper k -box of width at least m has a partition into proper boxes of width at least L_m , all having the original step and all axis lengths at most B_m .

Proof. Put $B = B_m$ and $q = \lceil n/B \rceil$. Partition the interval into q consecutive intervals with lengths $\lfloor n/q \rfloor$ and $\lceil n/q \rceil$. Their maximum length is at most B . Since $m \geq 4$, we have $B < m$, and hence $q \geq 2$. If $q = 2$, their minimum length is at least $\lfloor m/2 \rfloor = B - 1 \geq L_m$. If $q \geq 3$, then

$$\frac{n}{q} \geq \frac{(q-1)B + 1}{q} = B - \frac{B-1}{q} \geq \frac{2B+1}{3},$$

which again gives the required lower bound after taking floors. If m is even, $2B_m + 1 = m + 3$, and if m is odd, $2B_m + 1 = m + 2$; these identities show $L_m \geq m/3$ directly. The inequality $L_m \geq 2$ follows from $B_m \geq 3$. For a progression, apply the interval construction to its injective index parametrization. Apply this independently to the axes of a box and take the Cartesian products of the resulting parts. The index partitions give exact coverage and disjointness, and each part inherits properness and the original common step. $\square$

[write] 6 October 2026; sources 21–27. *Against the formal localization* (checked at HEAD 2dfe46f44). `Box.short_parent_partition` (`Proofs16ShortParents.lean`, line 23) localizes to width at least $m/8$, with axes of length at most half the modulus; the lemma here keeps $L_m \geq m/3$ with axes of length at most $\lfloor m/2 \rfloor + 1$. Not formalized.

Remark II.103.2 (Scope of the constant). The asymptotic constant $1/3$ is optimal for localization by consecutive index intervals subject to the cap B_m . Indeed, an axis of length $2B_m + 1$ needs at least three parts of length at most B_m ; one part therefore has length at most $\lfloor (2B_m + 1)/3 \rfloor = L_m$. This is a sharp statement about this localization problem. It is not an optimality claim for every conceivable proof of Lemma 16.10.

II.103.2 Transport with no further width loss

Lemma II.103.3 (Integer lift of a contained short progression). *Suppose $2(b-1) < p$. A proper progression $\{c + je : 0 \leq j < u\}$ of length $u \geq 2$, contained in the residues $\{0, \dots, b-1\}$, has a nonzero integer step a such that*

$$a \equiv e \pmod{p}, \quad |a|(u-1) \leq b-1.$$

The natural representatives, in their displayed order, form an ordinary integer progression of step a ; a may be negative.

Proof. Let $z_j \in \{0, \dots, b-1\}$ represent $c + je$. Every difference $z_{j+1} - z_j$ lies in $[-(b-1), b-1]$ and represents e modulo p . This interval contains at most one representative of any residue, because its length is less than p . All these differences are therefore one integer a . Properness gives $a \neq 0$, and the distance between the two endpoints is $|a|(u-1) \leq b-1$. $\square$

Lemma II.103.4 (Exact cylindrical transport). *Let R be a proper k -box with common step e and width at least the integer $w \geq 2$. Suppose one axis of R is contained in a proper progression $U = P(c, d, b)$, where $d \neq 0$, and suppose*

$$2(b-1) < p.$$

Let $J = P(z, d, n)$ be proper with $n \geq 2(b-1)$. Then $R \times J$ has a partition into proper $(k+1)$ -boxes, each with base exactly R , common step e , and width at least w .

More precisely, if the selected axis of R has length u , the number of parts can be taken to be an integer $t \geq 1$ satisfying

$$t(u-1) \leq b-1.$$

The final axes have lengths $\lfloor n/t \rfloor$ or $\lceil n/t \rceil$, and thus have length at least w .

Proof. Pull the selected axis back by the affine bijection $y \mapsto d^{-1}(y - c)$. Its carrier lies in $\{0, \dots, b-1\}$, it remains proper, and its step is $d^{-1}e$. By Lemma II.103.3, this step has a nonzero integer representative a with $|a|(u-1) \leq b-1$. Set $t = |a|$. Since $u \geq w \geq 2$,

$$tw \leq tu \leq 2t(u-1) \leq 2(b-1) \leq n.$$

Partition the index interval $\{0, \dots, n-1\}$ of J into its t residue classes modulo t . Each class is an integer progression of length $\lfloor n/t \rfloor$ or $\lceil n/t \rceil$, at least w . Use increasing order if $a > 0$ and decreasing order if $a < 0$. After the map $j \mapsto z + dj$, every final progression has step $da = e$ in $\mathbb{F}_p$. It is proper because it is a subsequence of the injective parametrization of J . These progressions partition J . Their products with R are consequently disjoint proper boxes, have the asserted width, and cover $R \times J$ exactly. $\square$

The integer-step and transport lemmas also hold in $\mathbb{Z}/N\mathbb{Z}$ for composite N , provided d is a unit and $2(b-1) < N$. Their proofs use no inverse of the child step. Primality is used in the subsequent interpolation theorem to invert distinct anchor differences, and in the common-partition construction to obtain a unit parent step from properness.

Remark II.103.5 (Why the length separation is real). The factor two in the sufficient condition $n \geq 2(b-1)$ cannot be decreased uniformly when children of width two are permitted. Let $b \geq 3$, choose a prime $p > 4b$, and take

$$U = \{0, \dots, b-1\}, \quad R = \{0, b-1\}, \quad J = \{0, \dots, 2b-4\}.$$

Here $|J| = 2(b-1) - 1$. Any box of width at least two contained in $R \times J$ has base R , and its common step must be $b-1$ or $-(b-1)$. The final coordinate $b-2$ has neither neighbor at one of these steps in J . Thus the point $(0, b-2)$ cannot belong to such a box. For example, $p = 17$, $R = \{0, 3\}$ and $J = \{0, 1, 2, 3, 4\}$ already give an obstruction at $(0, 2)$. Taking the Cartesian product of a refined base with the original final axis, without checking their steps, would overlook this obstruction.

II.103.3 Many slice covers on one base partition

Theorem II.103.6 (Common base refinement with exact transport). *Let $P = T \times J$ be a proper $(k+1)$ -box of width at least an integer $m \geq 4$. For $i = 1, \dots, r$, let $f_i : D_i \rightarrow \mathbb{F}_p$ be a partial function with $D_i \subseteq T$. Fix $a_i \in (0, 1]$, $\varepsilon_i \geq 0$, and positive integer Q_i . Assume the following cover statement for each i : for every proper k -box $U \subseteq T$, there is a partition of U into proper boxes of width at least $\text{width}(U)^{a_i}$, a set $G_{i,U} \subseteq U$ satisfying*

$$|U \setminus G_{i,U}| \leq \varepsilon_i |U|,$$

and, on each partition cell, at most Q_i multiaffine functions covering f_i on $D_i \cap G_{i,U}$ in that cell.

Put

$$A = \prod_{i=1}^r a_i, \quad \eta = \sum_{i=1}^r \varepsilon_i.$$

Then there are a common proper-box partition $\mathcal{R}$ of T and a set $G \subseteq T$ such that

$$|T \setminus G| \leq \eta |T|, \quad \text{width}(R) \geq L_m^A \geq (m/3)^A \quad (R \in \mathcal{R}).$$

On each $R \in \mathcal{R}$, every f_i is covered on $D_i \cap G \cap R$ by its own list of at most Q_i multiaffine functions. Furthermore, P has a partition into proper $(k+1)$ -boxes with these bases and width at least L_m^A . The product good set has $|(T \setminus G) \times J| \leq \eta |P|$.

Proof. First use Lemma II.103.1 to partition T into proper boxes T_α with widths at least L_m , original common step, and all axis lengths at most B_m . On each T_α , apply the cover assumption for f_1 . Apply that for f_2 to every resulting cell, and continue through f_r . This is a sequence of refinements, not an intersection of independently chosen partitions.

At stage i , the current cells partition T , so the union of their exceptional sets has size at most $\varepsilon_i |T|$. Let G be the complement of the union of all these stage exceptions. The union bound gives the claimed loss. Every final cell lies in a unique cell at each earlier stage, so the earlier lists restrict to that final cell without any increase of their lengths. This proves simultaneous coverage on G . Successively raising the width lower bounds to the powers $a_1, \dots, a_r$ gives L_m^A .

It remains to extend the base cells. The original common step d is nonzero because the input axes have length at least four. Each final base cell R lies in one of the T_α . Choose an axis of that parent, of length $b \leq B_m$, which contains the corresponding axis of R . Since $m \leq p$ and p is odd,

$$2(b-1) \leq 2 \lfloor m/2 \rfloor \leq p-1, \quad 2(b-1) \leq m \leq |J|.$$

Also $L_m \geq 2$ and $A > 0$, so $L_m^A > 1$; the integer width of R is therefore at least two. Apply Lemma II.103.4 with $w = \lceil L_m^A \rceil$. It partitions each $R \times J$ into genuine common-step boxes without decreasing the desired width. Different R have disjoint bases, so these partitions combine to partition P . Finally the product cardinality identity gives $|(T \setminus G) \times J| = |T \setminus G| |J|$. $\square$

[write] **6 October 2026; sources 21–27.** *A formalization candidate, not a formal result* (compared at HEAD 2dfe46f44). The formal local lift `section16_local_affine_lift_sqrt` (`Proofs16LiftWidth.lean`, line 31) has output width $\sqrt{(m/8)^A}/4$; the formal tiling `box_product_tiling_of_contained_axis` (`Proofs16RelativeSteps.lean`, line 89) needs $v^2 \leq \text{width} - 1$ and produces final axes of length $v - 1$ or v . The theorem here keeps width $L_m^A \geq (m/3)^A$ because it allows full residue chains as final axes, and it needs no inherited-coordinate hypothesis, which source 19's Corollary II.70.8 assumes. Whether the formal Section 16 interface (`Section16LineCover`, `Section16.lean`, line 691) accepts long final axes was not checked at the write; the source itself notes that long axes need not satisfy the Bohr budgets of Lemma 16.6 and argues that they are harmless at Lemma 16.10. No Lean statement of this theorem exists.

II.103.4 Affine interpolation and endpoint retention

Theorem II.103.7 (Geometric gluing of affine class covers). *Let $P = T \times J$ and $m \geq 4$ be as above. Let $f : D \rightarrow \mathbb{F}_p$, where $D \subseteq P$, and let $E \subseteq P$ satisfy $|P \setminus E| \leq \varepsilon_0 |P|$. For every $h \in T$, suppose the fibre*

$$\{x \in J : (h, x) \in D \cap E\}$$

is partitioned into at most q classes $C_t(h)$, on each of which $x \mapsto f(h, x)$ agrees with an affine function over $\mathbb{F}_p$.

Fix distinct anchors $x_1, \dots, x_r \in J$. Retain each complete class having at least two anchors, and retain every sampled point in each other class. Let D' denote the resulting subset of $D \cap E$, and assume

$$|(D \cap E) \setminus D'| \leq \delta |P|.$$

For the partial slice functions $f_i(h) = f(h, x_i)$, assume the cover hypotheses of Theorem II.103.6, with parameters a_i, ε_i, Q_i .

Then there are a set $H \subseteq P$ and a partition of P into proper boxes of width at least $L_m^A \geq (m/3)^A$, where $A = \prod_i a_i$, such that

$$|P \setminus H| \leq \left(\varepsilon_0 + \delta + \sum_i \varepsilon_i \right) |P|,$$

and on every partition cell the restriction of f to $D \cap H$ is covered by at most

$$M = \sum_i Q_i + \sum_{i < j} Q_i Q_j \tag{II.103.1}$$

multiaffine functions of all $k + 1$ coordinates. If D' instead retains only the complete classes with two anchors, the first sum in (II.103.1) can be omitted.

Proof. Construct the common base partition, its simultaneous good set G , and the full-box partition by Theorem II.103.6. On a base cell R , write its i th slice list as $\nu_{i,1}, \dots, \nu_{i,Q_i}$, padding if necessary. Include the functions $\nu_{i,u}(h)$, independent of x , and for $i < j$ include all functions

$$F_{i,j,u,v}(h, x) = \frac{(x - x_j)\nu_{i,u}(h) + (x_i - x)\nu_{j,v}(h)}{x_i - x_j}. \tag{II.103.2}$$

The denominator is nonzero, and hence invertible, because the anchors are distinct and the modulus is prime. Each displayed function is multiaffine in the base variables and in x ; the new variable only multiplies polynomials that are independent of it. The number of candidates is exactly the upper bound in (II.103.1) before deleting any duplications.

Set

$$H = E \cap (G \times J) \setminus ((D \cap E) \setminus D').$$

The union bound gives the asserted ambient loss. Consider $(h, x) \in D \cap H$ in one of the full-box cells. If $x = x_i$ is an anchor, the slice coverage gives $f(h, x) = \nu_{i,u}(h)$ for some u . Otherwise (h, x) belongs to a retained affine class with two anchors x_i, x_j . These two points lie in the domains of their slice functions, and $h \in G$, so their values are $\nu_{i,u}(h)$ and $\nu_{j,v}(h)$ for some u, v . The affine function on this class agrees with (II.103.2) at two distinct points. Subtracting the two affine functions shows their slopes and intercepts agree, so they agree at x as well.

The anchors need not lie in the final axis of the particular output cell. Both their slice covers and the affine identity refer to the original class in J ; restricting the candidate functions to a smaller final axis does not change that identity. This is why the transport step imposes no new sampling or coverage loss. $\square$

The following elementary sampling statement provides a completely uniform way to supply D' . It is included as a usable quantitative input, without a claim of a new optimal sampling constant.

Lemma II.103.8 (Finite-population sampling with retained observations). *Let J have $n \geq 1$ elements, and for every h in a nonempty finite set T let $C_1(h), \dots, C_q(h)$ be disjoint subsets of J . Choose a uniformly random r -element subset $S \subseteq J$, with $0 \leq r \leq n$. Retain every sampled point and every complete class meeting S in at least two elements. If $L(S)$ is the total number of discarded points, summed over h , then*

$$\frac{\mathbb{E}L(S)}{|T|n} \leq \left(1 - \frac{r}{n}\right) \min\left\{1, \frac{2q}{r+1}\right\}. \quad (\text{II.103.3})$$

In particular, for $0 < \delta \leq 1$ there is one shared sample with loss at most $\delta|T|n$ and

$$r = \min\{n, \lceil 2q/\delta \rceil\}.$$

Proof. The assertion is immediate for $r = n$, so suppose $r < n$. Fix h and then choose a uniform point y outside the random sample. Equivalently, choose a uniform $(r+1)$ -element set S^+ and then a uniform distinguished point $y \in S^+$, and put $S = S^+ \setminus \{y\}$. Write $K_t = |C_t(h) \cap S^+|$. The point y is discarded precisely when it belongs to a class for which $K_t \leq 2$. Thus the expected discarded mass in this fibre is

$$\frac{n-r}{r+1} \mathbb{E} \sum_t K_t \mathbf{1}_{\{K_t \leq 2\}}.$$

Disjointness gives $\sum_t K_t \leq r+1$, while there are at most q terms and each surviving term is at most two. The sum is at most $\min\{r+1, 2q\}$. Sum over h and divide by $|T|n$. An outcome no worse than the average supplies the shared sample. For the stated choice of r , either every point is sampled or $2q/(r+1) \leq \delta$. $\square$

[write] **6 October 2026; sources 21–27.** *Against source 19* (checked at the write). The add-one argument (choose $r+1$ points and mark one) is source 19's Lemma II.67.3, and its Theorem II.67.4 with interpolation order $d=2$ gives lost mass at most $\min\{1, 2q/(r+1)\}$ with $r = \min\{n, \lceil 2q/\sigma \rceil - 1\}$ anchors; source 19 also

has the exact minimax and the sharp constant (Theorems II.68.2, II.69.4). The bound here has the extra factor $1 - r/n$ because it counts a different loss: sampled points are retained as observations and paid for by the extra $\sum_i Q_i$ candidates of (II.103.1), while source 19 counts a sampled point of an unrecovered class as lost. Its stated sample size $\lceil 2q/\delta \rceil$ is one more than its own inequality needs and than source 19's at the same loss. In this report source 19's theorems are the first print and the sharper results; this lemma is printed as delivered. The formal `section16_distinct_sampling_budget` (Proofs16DistinctSampling.lean, line 117) uses $r \geq 6q/\sigma^2$.

In particular, short final axes require no separate pruning hypothesis: when $r = n$, the retained observations recover even singleton classes. The additional $\sum_i Q_i$ candidates in (II.103.1) pay for precisely these observations. For the remaining geometric input scales $m \in \{1, 2, 3\}$, partition the whole product into singleton boxes. On each such box the value of any partial function is covered by one constant function, with no loss and width $1 \geq (m/3)^A$. Thus the common-partition mechanism has an explicit small-scale branch as well as the case $m \geq 4$ proved above.

II.103.5 An explicit specialization to the Section 16 controls

Corollary II.103.9 (Equal slice profiles). *Assume the hypotheses of Theorem II.103.7, and suppose there is at least one anchor ($r \geq 1$) and every sampled slice has the following common profile: an exceptional fraction ε permits a width exponent*

$$a(\varepsilon) = \left(\frac{\gamma\varepsilon}{s} \right)^{D_k s}, \quad Q(\varepsilon) = a(\varepsilon)^{-1}, \quad D_k = 2^{2^{k+8}},$$

with $0 < \gamma \leq 1$ and $s \geq 1$. Here an integer list length is bounded by the displayed real number $Q(\varepsilon)$. Give the r slices a total exceptional fraction $0 < \eta \leq 1$, using $\varepsilon = \eta/r$ for each. Then the final width is at least

$$L_m^{A_*} \geq (m/3)^{A_*}, \quad A_* = \left(\frac{\gamma\eta}{rs} \right)^{D_k r s},$$

and the number of candidates on each cell is at most

$$rQ_* + \binom{r}{2} Q_*^2, \quad Q_* = \left(\frac{rs}{\gamma\eta} \right)^{D_k s}.$$

The ambient loss is at most $\varepsilon_0 + \delta + \eta$.

Proof. Substitute the slice profile into Theorem II.103.7. The product of the r identical width exponents is A_* . □

There is also a profile uniform in the size of the final axis. Assume the displayed slice-cover profile is available for every possible anchor in J , before choosing the sample. Given the fibre-class bound $q \geq 1$ and a loss budget $0 < \delta \leq 1$, put $R_0 = \lceil 2q/\delta \rceil$. Lemma II.103.8 uses $r = \min\{|J|, R_0\}$ distinct anchors. Give each actual slice the budget η/R_0 , so the total slice loss is at most η . Then all final boxes have width at least

$$L_m^{A_0}, \quad A_0 = \left(\frac{\gamma\eta}{R_0 s} \right)^{D_k R_0 s},$$

and have at most

$$R_0 Q_0 + \binom{R_0}{2} Q_0^2, \quad Q_0 = \left(\frac{R_0 s}{\gamma\eta} \right)^{D_k s}$$

candidates. Indeed the actual exponent is $a^r \geq a^{R_0} = A_0$ for $a = Q_0^{-1} \in (0, 1]$, and the candidate count increases with the number of slices. This gives one bound for both $|J| \leq R_0$ and $|J| > R_0$, without padding a short axis with nonexistent distinct anchors.

For unequal parameters $s_i \geq 1$, with a common γ and D_k , the analogous product of width exponents is maximized, subject to $\sum_i \varepsilon_i = \eta$, by

$$\varepsilon_i = \eta \frac{s_i}{S}, \quad S = \sum_i s_i.$$

Indeed its logarithm is $D_k \sum_i s_i \log(\gamma \varepsilon_i / s_i)$, a strictly concave function; its constrained critical point satisfies $s_i / \varepsilon_i = S / \eta$. The optimized exponent is $(\gamma \eta / S)^{D_k S}$, with the exact heterogeneous candidate count still given by (II.103.1).

II.103.6 Integration boundary and further questions

The conditional inputs above match the two types of data appearing in the Lemma 16.10 argument: affine fibre lists from Lemma 16.9 and multiply multilinear slice covers from the induction hypothesis. The construction supplies an actual common base partition, all needed step alignment, exact coverage of the resulting product partition, and a small-final-axis sampling branch. The comparison is specifically with the proper-box formulation used at repository revision 65f4e039fc4ad30a8f492e63abde2bcd0740d7c9.

The existing modules `Proofs16Interpolation.lean` and `Proofs16RelativeSteps.lean` already contain the two-anchor algebra and related transport tools. The latter's named product tiling theorem imposes $v^2 \leq \text{width}(R) - 1$ and produces final lengths $v - 1$ or v . Our transport theorem has a different purpose: it permits full residue progressions as final axes and thus preserves the child width. Those long axes would not automatically satisfy the Bohr-linearity budgets needed earlier in Lemma 16.6. They are legitimate here because the relevant functions have already been shown affine on each original fibre class. No new Lean verification is asserted by the present ordinary mathematical proof.

These results also do not prove the displayed final parameter comparison in the repository's Lemma 16.10. The exponent A_* still depends on the number r of sampled slices, which in turn depends on the fibre-cover count and the requested loss. Substituting these dependencies and comparing them with the claimed $c(\rho, \gamma, k + 1)$ is a separate quantitative obligation. Nor does a result for proper modular boxes by itself settle the repository's separate question about nonwrapping progressions in integer representatives. The argument above explicitly uses modular boxes and local affine index pullbacks.

II.104 Sources 31 and 29: the Proposition 6.1 endpoint [write]

Added in the sixth write, 6 October 2026.

The next nine sections are about the endpoint of Gowers's Proposition 6.1: functions whose derivative spectrum is almost as concentrated as possible, and selectors of derivative frequencies that capture almost all of it. *Logically the chapter continues source 02's Sections II.2–II.4 and source 11's chapter* (Section II.22, with its Corollary II.32.2); it is printed here, at the end of Part II, because no number printed earlier may move. Sections II.105–II.108 and II.109 are source 31, the main source; Sections II.110–II.113 are source 29's, written the same day from the same pin, a

second route with cruder constants and with soft and mixed selectors. Neither cites the other. Their introductions are Sections F.36 and F.34; source 31's Sections 6–8 form the next chapter. The notation is in the table at the head of this Part.

What is proved, in one paragraph. For every finite abelian group and nonzero f , with $\kappa = \|f\|_{U^3}^8 / \|f\|_2^8$ and $\delta = 1 - \kappa < 1/120$, the best squared correlation F with a normalized quadratic phase on a coset satisfies $1 - F \leq \mathcal{R}(\delta) = (1 - \sqrt{(1 + \sqrt{1 - 4\delta})/2})/2$, and this function is optimal, with equality on $\mathbb{Z}/2\mathbb{Z}$ (source 31's Theorem II.107.4; the preliminary bound 20δ is Theorem II.106.6, the quartic envelope Theorem II.107.3). For a selected derivative graph of score $1 - \varepsilon$, $\varepsilon \leq 1/400$, the graph errors are at most ε and the infidelity at most $(1 - \sqrt{1 - 2\varepsilon})/2$, both separately sharp (Theorem II.108.1); a weight-defect inequality controls the weights in the U^2 and coset topology, and explicit examples show that near equality cannot control them in normalized L^1 without complexity assumptions (Section II.109). Source 29 proves the same endpoint with graph error 20ε and squared distance 22ε for $\varepsilon \leq 2^{-9}$ (Theorem II.110.1), and extends it to soft selectors and to two functions (Theorems II.111.1, II.112.1); the linear order is sharp (Section II.113). The qualitative near-extremizer classification is Eisner and Tao's; both sources credit it.

Relation to the formal statements (checked at HEAD 5d4718772). `proposition_6_1` (Section ns06_07.lean, line 37) has the exact companion `proposition_6_1_holds` (Proofs05_10.lean, line 505). No near-extremal stability statement is formal; nothing in this chapter is formalized.

II.105 Normalized derivatives and their phase space

Source 31, delivered Section 2.

Let G be a finite abelian group, written additively, and let $N = |G|$. Spatial averages are normalized. Fourier sums over the character group $\widehat{G}$ use counting measure. We use the inner product $\langle u, v \rangle = \mathbb{E}_x u(x) \overline{v(x)}$, linear in its first argument. For $f \neq 0$, set

$$m = \mathbb{E}_x |f(x)|^2, \quad v = f/\sqrt{m}, \quad \Delta_h f(x) = f(x) \overline{f(x-h)}.$$

Write

$$A_f(h, \chi) = \mathbb{E}_x \Delta_h f(x) \overline{\chi(x)}, \tag{II.105.1}$$

$$q_f(h, \chi) = \frac{|A_f(h, \chi)|^2}{m^2}, \tag{II.105.2}$$

$$\kappa(f) = \frac{\|f\|_{U^3}^8}{m^4}. \tag{II.105.3}$$

The U^3 norm uses the usual alternating-conjugation cube average. Equivalently,

$$\|f\|_{U^3}^8 = \mathbb{E}_h \sum_{\chi \in \widehat{G}} |A_f(h, \chi)|^4. \tag{II.105.4}$$

This follows from the identity $\|g\|_{U^2}^4 = \sum_{\chi} |\widehat{g}(\chi)|^4$ applied to each derivative. All these expressions include zero directions and repeated cube vertices.

The phase space is the abelian group $V = G \times \widehat{G}$, of size N^2 . For $z = (h, \chi)$ define the unitary Weyl operator

$$W_z = M_\chi T_h, \quad (M_\chi u)(x) = \chi(x)u(x), \quad (T_h u)(x) = u(x-h).$$

Then $q_f(z) = |\langle v, W_z v \rangle|^2$. Multiplication of Weyl operators agrees with addition in V up to a scalar of modulus one. Their commutator is the alternating bicharacter

$$B((h, \chi), (k, \psi)) = \chi(k) \overline{\psi(h)}. \quad (\text{II.105.5})$$

Thus $W_z W_{z'} = B(z, z') W_{z'} W_z$. This convention fixes every sign used below; changing all signs consistently leaves the squared coefficients unchanged.

Lemma II.105.1 (Mass, critical norm, and self-duality). *For $q = q_f$ one has*

$$0 \leq q(z) \leq 1, \quad \sum_{z \in V} q(z) = N, \quad \kappa = \frac{1}{N} \sum_{z \in V} q(z)^2 \leq 1. \quad (\text{II.105.6})$$

Define the unitary symplectic Fourier transform on V by

$$(\mathcal{F}a)(z) = \frac{1}{N} \sum_{y \in V} \overline{B(z, y)} a(y).$$

Then $\mathcal{F}q = q$.

Proof. The pointwise bound is Cauchy–Schwarz. Parseval in χ gives

$$\sum_{h, \chi} |A_f(h, \chi)|^2 = \sum_h \mathbb{E}_x |f(x)|^2 |f(x - h)|^2 = Nm^2.$$

Equation (II.105.4) gives the third identity, and $q^2 \leq q$ gives the inequality. In particular this proves the finite critical inequality $\|f\|_{U^3} \leq \|f\|_2$ directly.

For self-duality one can expand the squared coefficient. At $z = (a, \alpha)$, summation over χ forces the two base variables to differ by a . With $m = 1$ the remaining expression is

$$\frac{1}{N^2} \sum_{h, x} f(x) \overline{f(x - h)} \overline{f(x - a)} f(x - a - h) \overline{\alpha(h)}.$$

Substitute $y = x - h$. The two sums separate into a quantity and its conjugate, giving $|A_f(a, \alpha)|^2 = q_f(a, \alpha)$. Scaling handles general m . Character orthogonality proves unitarity of $\mathcal{F}$; its kernel satisfies $B(y, z) = \overline{B(z, y)}$, so it is also self-adjoint. $\square$

II.105.1 The weighted derivative-energy inequality

For a nonnegative weight $w : V \rightarrow \mathbb{R}_{\geq 0}$ define

$$\Theta_w(f) = \frac{1}{N} \sum_{z \in V} w(z) q_f(z), \quad (\text{II.105.7})$$

$$\mathbf{E}(w) = \sum_{z_1 + z_2 = z_3 + z_4} w(z_1) w(z_2) w(z_3) w(z_4), \quad (\text{II.105.8})$$

$$\mathcal{E}(w) = N^{-3} \mathbf{E}(w). \quad (\text{II.105.9})$$

Notice that the energy is divided by N^3 , not by $|V|^3 = N^6$. The distinction matters because the extremal subgroups of V have size N . When $0 \leq w \leq 1$, the score Θ_w lies in $[0, 1]$.

Proposition II.105.2 (Weighted energy certificate). *For every such w ,*

$$\Theta_w(f)^4 \leq \kappa(f)\mathcal{E}(w). \quad (\text{II.105.10})$$

Proof. Self-duality and Parseval give $\Theta_w = N^{-1} \sum_z q(z) \overline{(\mathcal{F}w)(z)}$. Apply Cauchy–Schwarz first with weights q/N , whose sum is one, and then with counting measure:

$$\Theta_w^2 \leq \frac{1}{N} \sum_z q(z) |\mathcal{F}w(z)|^2, \quad \Theta_w^4 \leq \left(\frac{1}{N} \sum_z q(z)^2 \right) \left(\frac{1}{N} \sum_z |\mathcal{F}w(z)|^4 \right).$$

Fourier orthogonality on the group of size N^2 gives $\sum_z |\mathcal{F}w(z)|^4 = N^{-2} \mathbf{E}(w)$. The desired inequality follows. $\square$

This is the single-function specialization of the weighted refinement of Gowers’ Proposition 6.1 in the existing project report [140]. We include a short phase-space proof to make the later extraction independent of normalization conventions. The proposition itself is a credited prerequisite, not an additional novelty claim.

Quantity	This article	Unnormalized spatial Fourier sum
Derivative coefficient	$A_f(h, \chi)$	$NA_f(h, \chi)$
Selected graph score	$m^2 \Theta_{\mathbf{1}_\Gamma}$	$N^3 m^2 \Theta_{\mathbf{1}_\Gamma}$
Additive graph energy	$\mathbf{E}(\mathbf{1}_\Gamma)$	Same ordered quadruple count
Critical ratio	$\kappa = \ f\ _{U_3}^8 / m^4$	Dimensionless

II.106 Recovering an extremal subgroup without an exponent loss

Source 31, delivered Section 3.

Definition II.106.1 (Quadratic coset phases). A unit phase $Q : K \rightarrow \mathbb{T}$ on a subgroup $K \leq G$ is quadratic if all its third multiplicative derivatives are one. Equivalently, each derivative has the form $\Delta_h Q(x) = c_h \chi_h(x)$ with $\chi_h \in \widehat{K}$, where $h \mapsto \chi_h$ is a symmetric homomorphism. A normalized quadratic coset phase on G is

$$\psi(x) = \left(\frac{N}{|K|} \right)^{1/2} \mathbf{1}_{x_0+K}(x) Q(x - x_0). \quad (\text{II.106.1})$$

Let $\mathcal{Q}_{\text{cos}}(G)$ denote this collection, including arbitrary constant unit phases. Its elements have L^2 norm one. Define the fidelity

$$F(f) = \max_{\psi \in \mathcal{Q}_{\text{cos}}(G)} |\langle f / \sqrt{m}, \psi \rangle|^2. \quad (\text{II.106.2})$$

The maximum exists: apart from global phase, there are finitely many subgroups, cosets, symmetric bicharacters, and quadratic refinements on a finite group. We will also identify these phases as common eigenvectors of maximal commuting Weyl families, which gives another finite description.

II.106.1 Two elementary closure lemmas

Lemma II.106.2 (Projective overlap). *If $q(z), q(z') \geq 1 - a$, with $0 \leq a \leq 1/2$, then*

$$q(z - z') \geq (1 - 2a)^2 \geq 1 - 4a.$$

Proof. For unit vectors u_1, u_2 each having squared overlap at least $1 - a$ with the same unit vector v , orthogonal decomposition along v gives

$$|\langle u_1, u_2 \rangle| \geq (1 - a) - a = 1 - 2a.$$

Apply this to suitable scalar multiples of $W_z v$ and $W_{z'} v$. Their overlap has squared modulus $q(z - z')$, because Weyl products differ from the operator with the summed label only by a unit scalar. $\square$

Lemma II.106.3 (Small difference set). *If a finite nonempty subset S of an abelian group satisfies $|S - S| < 3|S|/2$, then $S - S$ is a subgroup.*

Proof. Put $D = S - S$. For $d = a - b \in D$, the sets $S - a$ and $S - b$ both lie in D , so their intersection has size greater than $|S|/2$. Consequently $|S \cap (S + d)| > |S|/2$. For $d, e \in D$, the subsets of S whose translates by d and by e remain in S must intersect. Thus $d - e \in D$. Since D is nonempty and closed under subtraction, it is a subgroup. $\square$

II.106.2 Averaging forces exact commutativity

A subgroup $H \leq V$ is called isotropic when $B(h, h') = 1$ for all $h, h' \in H$. The following argument is useful because its lower gap is $1/2$, independent of the orders of the roots of unity that occur.

Lemma II.106.4 (Noncommuting subgroup gap). *If $H \leq V$ is not isotropic, then for every unit vector v ,*

$$\frac{1}{|H|} \sum_{h \in H} |\langle v, W_h v \rangle|^2 \leq \frac{1}{2}. \quad (\text{II.106.3})$$

Proof. Let P be the rank-one orthogonal projection onto v , and put

$$\rho = \frac{1}{|H|} \sum_{h \in H} W_h P W_h^*.$$

Conjugation is an honest group action even though the Weyl operators themselves form a projective representation. Therefore ρ commutes with every W_h . It is positive, has trace one, and the averaging map is an orthogonal projection for the Hilbert–Schmidt inner product. In particular,

$$\frac{1}{|H|} \sum_h |\langle v, W_h v \rangle|^2 = \text{tr}(P\rho) = \text{tr}(\rho^2).$$

Choose $a, b \in H$ with $W_a W_b = \zeta W_b W_a$, where $\zeta \neq 1$ has order $r \geq 2$. Every eigenspace of ρ is invariant under both operators. Within a nonzero eigenspace, choose an eigenvector of W_a . Its images under $1, W_b, \dots, W_b^{r-1}$ have distinct W_a eigenvalues and are linearly independent. Hence every nonzero eigenvalue of ρ has multiplicity at least r , and is at most $1/r$. It follows that $\text{tr}(\rho^2) \leq (1/r) \text{tr} \rho \leq 1/2$. $\square$

[write] **6 October 2026; sources 34–39.** Source 36’s Theorem IV.17.1 is the sharp, rank-sensitive form of this lemma over $\mathbb{F}_p^n$ (bound p^{-r} averaged over half-ranks), and source 38 reproves the lemma for odd order (its Lemma II.150.3); neither source 36 nor this source cites the other.

Lemma II.106.5 (Maximal isotropic subgroups and coset phases). *If $H \leq V$ is isotropic, then $|H|$ divides N . If $|H| = N$, its common Weyl eigenspaces are one dimensional and their unit vectors are quadratic coset phases. Conversely every quadratic coset phase has*

$$q_\psi = \mathbf{1}_H \quad (\text{II.106.4})$$

for an isotropic subgroup H of size N .

Proof. The operators W_h , $h \in H$, commute and are unitary, so they are simultaneously diagonalizable. Choose a common eigenvector with eigenvalues λ_h , and divide W_h by λ_h . The resulting operators U_h form an honest representation of H : the projective multiplication scalars cancel by evaluating them on the chosen eigenvector. For $\gamma \in \hat{H}$, the character projector is

$$P_\gamma = \frac{1}{|H|} \sum_{h \in H} \overline{\gamma(h)} U_h.$$

The trace of W_h is zero for $h \neq 0$: a nonzero translation has no diagonal entries, and a nontrivial pure modulation has character sum zero. Thus $\text{tr } P_\gamma = N/|H|$. A projector has integer trace, proving divisibility; when $|H| = N$, each projector has rank one.

For the structural description let K be the projection of H to G , and let $D = \{\chi : (0, \chi) \in H\}$. Isotropy implies $D \subseteq K^\perp$. When $|H| = N$, cardinalities give $D = K^\perp$. The modulation eigenrelations force the support of a common eigenvector into one coset $x_0 + K$. The translation eigenrelations for K force every point of that coset into the support and make the modulus constant there. Normalization gives the amplitude in (II.106.1). The same eigenrelations say that every derivative on K is a scalar times a character. Their commutativity supplies symmetry, and hence vanishing third multiplicative derivatives.

Conversely take (II.106.1). For each $h \in K$, its derivative character on K has exactly $|K^\perp|$ extensions to G . Together these extensions give a subgroup $H \leq V$ of size $|K||K^\perp| = N$. Its Weyl operators preserve the line through ψ and therefore commute on that line; their scalar commutators are consequently one. Thus H is isotropic and $q_\psi = 1$ on H . The mass identity (II.105.6) forces $q_\psi = 0$ elsewhere. $\square$

Theorem II.106.6 (A uniform preliminary stability bound). *If $\delta = 1 - \kappa(f) < 1/120$, then there is an isotropic subgroup $H \leq V$ of size N , and a common eigenstate $\psi \in \mathcal{Q}_{\text{cos}}(G)$, such that*

$$|\langle v, \psi \rangle|^2 \geq 1 - 20\delta. \quad (\text{II.106.5})$$

One can take $H = S - S$ with $S = \{z : q_f(z) \geq 19/20\}$.

Proof. Since $\sum_z q(z)(1 - q(z)) = N\delta$, the q -mass outside S is at most $20N\delta$. Therefore

$$|S| \geq \sum_{z \in S} q(z) \geq (1 - 20\delta)N.$$

Lemma II.106.2 gives $S - S \subseteq \{q \geq 4/5\}$, so $|S - S| \leq 5N/4$. The strict inequality $\delta < 1/120$ makes $5N/4 < 3|S|/2$. Lemma II.106.3 therefore makes $H = S - S$ a subgroup. Every coefficient on H is

at least $4/5$; Lemma II.106.4 makes H isotropic. Also $|H| \geq |S| > 5N/6 > N/2$, so divisibility in Lemma II.106.5 gives $|H| = N$.

We have $0 \in S$, hence $S \subseteq H$. Expand v in the common eigenbasis of H , and let the squared coefficients be p_γ . Character orthogonality gives

$$\frac{1}{N} \sum_{h \in H} q(h) = \sum_{\gamma} p_\gamma^2 \geq 1 - 20\delta. \quad (\text{II.106.6})$$

The p_γ sum to one, so their maximum is at least their sum of squares. A maximizing basis vector is the required ψ . $\square$

The particular threshold $19/20$ is a convenient rational choice. Optimizing it would slightly enlarge the preliminary range; the final sharp modulus below comes from a different calculation and is not limited by the coefficient 20 as a local error constant.

II.107 The sharp local quartic envelope

Source 31, delivered Section 4.

Fix a maximal isotropic subgroup H and a common unit eigenvector ψ . Assume $|\langle v, \psi \rangle|^2 = F = 1 - t$, where $0 \leq t \leq 1/6$. Index the common eigenbasis so that ψ has label zero. Write its coefficient probabilities as

$$p_0 = F, \quad r_\gamma = p_\gamma \ (\gamma \neq 0), \quad r_0 = 0, \quad \sum_{\gamma} r_\gamma = t.$$

The labels form the abelian group $\hat{H}$. Define

$$A = \sum_{\gamma} r_\gamma^2, \quad B_0 = \sum_{\gamma} r_\gamma r_{-\gamma}, \quad C_0 = \sum_{a+b=c} r_a r_b r_c. \quad (\text{II.107.1})$$

We distinguish B_0 from the symplectic bicharacter B .

Lemma II.107.1 (Tail estimates). *These statistics satisfy*

$$B_0 \leq A \leq t^2, \quad C_0 \leq \frac{3}{2}t(t^2 - A), \quad \mathbb{E}(r) \leq t^2 A.$$

If H has odd order, then also $B_0 \leq t^2 - A$.

Proof. The first assertions are Cauchy–Schwarz and nonnegativity. In a relation $a + b = c$ with $r_a r_b r_c > 0$, all three labels are nonzero; they cannot all coincide. At least two of the three pairwise comparisons are unequal. Summing this observation with weights gives $2C_0 \leq 3t(t^2 - A)$. For the energy bound, Young’s inequality gives $\|r * r\|_2 \leq \|r\|_1 \|r\|_2 = t\sqrt{A}$; alternatively this follows by applying Cauchy–Schwarz to each convolution sum and then summing. Finally, in odd order every nonzero label differs from its negative, so the terms defining B_0 occur among the off-diagonal terms whose total is $t^2 - A$. $\square$

Lemma II.107.2 (Outside coefficients). *For $z \notin H$,*

$$q_f(z) \leq 4Ft. \quad (\text{II.107.2})$$

Proof. Maximality and nondegeneracy of the symplectic pairing imply $H^\perp = H$. Therefore $W_z\psi$ and $W_z^*\psi$ are common eigenvectors with labels different from that of ψ , and are orthogonal to it. It follows that $s = |\langle W_z v, \psi \rangle|^2 \leq t$. Decompose v and $W_z v$ along the line through ψ and its orthogonal complement. Cauchy–Schwarz gives

$$|\langle v, W_z v \rangle| \leq \sqrt{Fs} + \sqrt{t(1-s)}.$$

For $0 \leq s \leq t \leq F$ the right side is increasing in s . Its maximum is $2\sqrt{Ft}$, proving the assertion. $\square$

Theorem II.107.3 (Quartic envelope and its deficit). *Under the preceding hypotheses,*

$$\kappa(f) \leq \Phi(t) := (1-t)^4 + 14(1-t)^2 t^2 + t^4 = 1 - 4t + 20t^2 - 32t^3 + 16t^4. \quad (\text{II.107.3})$$

More precisely, with $a(t) = 6(1-t)^2 - 10(1-t)t + t^2$,

$$\Phi(t) - \kappa(f) \geq a(t)(t^2 - A) + 2(1-t)^2(A - B_0). \quad (\text{II.107.4})$$

For odd $|G|$ this implies the following bound, strictly smaller when $t > 0$:

$$\kappa(f) \leq (1-t)^4 + 12(1-t)^2 t^2 + t^4. \quad (\text{II.107.5})$$

The universal envelope (II.107.3) is attained on $G = \mathbb{Z}/2\mathbb{Z}$.

Proof. On H the Weyl coefficient is, up to a unit scalar, the Fourier transform of p . Its fourth moment is the additive energy of p . Classify an additive quadruple by the positions occupied by nonzero labels. One nonzero position contributes nothing; two opposite-side positions contribute $4F^2A$; two same-side positions contribute $2F^2B_0$; three positions contribute $4FC_0$. Hence

$$\frac{1}{N} \sum_{h \in H} q(h)^2 = F^4 + 4F^2A + 2F^2B_0 + 4FC_0 + \mathbf{E}(r). \quad (\text{II.107.6})$$

This identity counts positions, so repetitions are already included. Equation (II.106.6) without its inequality says that the q -mass outside H , divided by N , is $1 - F^2 - A$. Thus Lemma II.107.2 bounds its fourth-moment contribution by $4Ft(1 - F^2 - A)$. Apply Lemma II.107.1 to obtain

$$\kappa \leq F^4 + 6Ft^3 + 4Ft(1 - F^2) + a(t)A - 2F^2(A - B_0).$$

Replacing A by t^2 in the first occurrence gives exactly $F^4 + 14F^2t^2 + t^4$. Since $a(t) > 0$ for $t \leq 1/6$, this proves (II.107.3) and (II.107.4).

For odd order, put $x = t^2 - A$. Then $A - B_0 \geq \max(0, t^2 - 2x)$. Also

$$a(t) - 4F^2 = 2 - 14t + 13t^2 \geq 1/36 > 0 \quad (0 \leq t \leq 1/6).$$

Minimizing $a(t)x + 2F^2 \max(0, t^2 - 2x)$ over $x \geq 0$ gives $2F^2t^2$. Subtract this from (II.107.3) to obtain (II.107.5). The equality example is calculated below. $\square$

[write] **6 October 2026; sources 34–39.** The optimal odd-order coefficients, which source 31 leaves open (its Research questions), are source 38’s Theorem II.150.11: $\varepsilon/4 + 3\varepsilon^2/16 + (\sqrt{2}/64)\varepsilon^{5/2} + O(\varepsilon^3)$, the fractional term from characters of order three (item II.46).

II.107.1 Exact inversion and optimality

Theorem II.107.4 (Sharp local critical-norm stability). *Let G be any finite abelian group and $f \neq 0$. If $\delta = 1 - \kappa(f) < 1/120$, then*

$$1 - F(f) \leq \mathcal{R}(\delta) = \frac{1 - \sqrt{(1 + \sqrt{1 - 4\delta})/2}}{2}. \quad (\text{II.107.7})$$

Equivalently,

$$\min_{\psi \in \mathcal{Q}_{\cos}(G)} \|f/\sqrt{m} - \psi\|_2^2 \leq 2(1 - \sqrt{1 - \mathcal{R}(\delta)}). \quad (\text{II.107.8})$$

For every $0 \leq \delta < 1/120$, equality is possible in both formulas. Thus the functions on their right sides are optimal uniformly over finite abelian groups on this interval.

Proof. Theorem II.106.6 gives $F(f) \geq 1 - 20\delta$. Choose a maximizing coset phase in (II.106.2), and its maximal isotropic subgroup from Lemma II.106.5. Then $t = 1 - F(f) < 1/6$, so Theorem II.107.3 applies:

$$\delta \geq P(t) := 1 - \Phi(t) = 4t(1 - t)(1 - 2t)^2. \quad (\text{II.107.9})$$

On $[0, 1/6]$ one has $P(t) \geq (40/27)t$, so $t \leq 27\delta/40 < 9/1600$. In particular t lies in the increasing branch of P , which extends to $(2 - \sqrt{2})/4$.

Set $x = 4t(1 - t)$. Then $P(t) = x(1 - x)$. Inverting first this quadratic and then $x = 4t(1 - t)$ on the smaller branches gives (II.107.7). Optimizing the constant phase of ψ gives the identity $\min_{\psi} \|v - \psi\|_2^2 = 2(1 - \sqrt{F(f)})$, which proves (II.107.8).

For equality, take $G = \mathbb{Z}/2\mathbb{Z}$ and, in the unit point-mass basis for normalized L^2 , let

$$v = \sqrt{1 - t} e_0 + \sqrt{t} e_1. \quad (\text{II.107.10})$$

The four squared Weyl coefficients are

$$1, \quad (1 - 2t)^2, \quad 4t(1 - t), \quad 0.$$

Their squared sum divided by two is $\Phi(t)$. The coset phases on this group are the two point masses and, up to global phase, $(e_0 + \zeta e_1)/\sqrt{2}$ for $\zeta \in \{1, -1, i, -i\}$. For $t < (2 - \sqrt{2})/4$, their greatest fidelity with (II.107.10) is $1 - t$. Taking $t = \mathcal{R}(\delta)$ proves simultaneous equality. $\square$

[write] **6 October 2026; sources 28–33.** *The strongest endpoint result in this report* (checked at the write with the intake's enumeration: on $\mathbb{Z}/p\mathbb{Z}$, $p \leq 13$, 764 cases with $\delta < 1/120$, $\max(1 - F)/\mathcal{R}(\delta) = 0.999997$). It supersedes source 29's constants for the same endpoint (note after its Theorem II.110.1) and answers source 02's question on near equality for the function in a sharp form (item II.2); the qualitative statement is Eisner and Tao's, as the source says.

[write] **7 October 2026; sources 54–59.** *The second coefficient in every degree* (source 55, Section II.236): the squared distance on the right of (II.107.8) is $\delta/4 + 21\delta^2/64 + O(\delta^3)$, and source 55's Theorem II.237.1 gives, for $|f| \leq 1$ and phase polynomials of degree at most $d - 1$, the coefficients 2^{1-d} and $(q^2 - 1)/(3q^2)$ on even-order groups ($21/64$ at $d = 3$), $(3^d - 1)/(2q^2)$ on odd-order groups, both sharp (note there). Source 55 does not cite this source.

The local expansion is explicit:

$$\mathcal{R}(\delta) = \frac{\delta}{4} + \frac{5\delta^2}{16} + \frac{21\delta^3}{32} + \frac{429\delta^4}{256} + O(\delta^5), \quad (\text{II.107.11})$$

$$2(1 - \sqrt{1 - \mathcal{R}(\delta)}) = \frac{\delta}{4} + \frac{21\delta^2}{64} + \frac{357\delta^3}{512} + \frac{29325\delta^4}{16384} + O(\delta^5). \quad (\text{II.107.12})$$

Thus the sharp first-order constant for squared L^2 distance is $1/4$. The error variable is $1 - \|f\|_{U^3}^8/\|f\|_2^8$, not $1 - \|f\|_{U^3}/\|f\|_2$; their first-order conversion differs by a factor of eight.

II.107.2 What near equality in the envelope forces

The deficit formula supplies more than an extremal value. If $t > 0$ and $\Delta = \Phi(t) - \kappa$, then

$$\max_{\gamma \neq 0} \frac{r_\gamma}{t} \geq \frac{A}{t^2} \geq 1 - \frac{\Delta}{a(t)t^2}.$$

So a near-extremizer at a fixed small infidelity has almost all its remaining mass in a single other common eigenvector. At exact equality, $A = t^2$ and $B_0 = A$ force that other label to have order two. The outside-coefficient bound must also be saturated wherever the outside coefficient is nonzero. Conversely, these conditions yield the universal envelope: a single order-two tail label makes $C_0 = 0$, $E(r) = t^4$, and the saturated outside contribution is $8F^2t^2$.

This describes the two-state mechanism behind sharpness and explains the strict improvement in odd order. Formula (II.107.5) is only an upper envelope; determining the optimal odd-order second coefficient remains an interesting separate problem.

II.108 Sharp stability of a selected derivative graph

Source 31, delivered Section 5.

Let $B \subseteq G$ and let $\phi : B \rightarrow \widehat{G}$. Its normalized score is

$$T_\phi = \frac{1}{N} \sum_{h \in B} q_f(h, \phi(h)).$$

For the full-domain case $B = G$, this is the square-correlation form of the input used at Gowers' Proposition 6.1. Partial graphs are allowed with the same ambient normalization.

Theorem II.108.1 (Phase error, graph error, and missing directions). *Suppose $T_\phi \geq 1 - \varepsilon$, where $0 \leq \varepsilon \leq 1/400$. There is a unit quadratic phase Q on all of G , with symmetric derivative homomorphism $L : G \rightarrow \widehat{G}$, such that, with*

$$t = 1 - |\langle v, Q \rangle|^2, \quad z = 1 - |B|/N, \quad \beta = N^{-1}|\{h \in B : \phi(h) \neq L(h)\}|,$$

one has

$$\varepsilon \geq 2t(1 - t) + z(1 - 4t(1 - t)) + \beta(1 - 8t(1 - t)). \quad (\text{II.108.1})$$

In particular,

$$t \leq \mathcal{G}(\varepsilon) := \frac{1 - \sqrt{1 - 2\varepsilon}}{2}, \quad z + \beta \leq \varepsilon. \quad (\text{II.108.2})$$

The best phase is attained within the full-support quadratic phases, and L is the unique homomorphism having at most εN disagreements on B when missing directions are also charged. The two bounds in (II.108.2) are separately sharp.

Proof. Cauchy–Schwarz on the N possible directions, putting zero on missing ones, gives $\kappa \geq T_\phi^2 \geq (1 - \varepsilon)^2$. Thus $\delta \leq 2\varepsilon - \varepsilon^2 < 1/120$. Choose a maximizing coset phase Q by Theorem II.107.4, and let H be its extremal subgroup. The preliminary estimate already gives $t \leq 20\delta \leq 1/10$.

If the projection K of H into G were proper, it would have density at most $1/2$. For directions outside K , every selected coefficient is outside H , so Lemma II.107.2 gives

$$T_\phi \leq \frac{1}{2} + 2t(1-t) \leq \frac{1}{2} + \frac{1}{5} < 1 - \varepsilon.$$

This contradiction forces $K = G$. Consequently H is the graph of a homomorphism $L : G \rightarrow \widehat{G}$, and isotropy makes it symmetric. Lemma II.106.5 makes Q a full-support unit phase.

Put $u = t(1-t)$. On H the coefficient amplitude is the Fourier transform of the probability vector with atom $1-t$ at zero and remaining mass t . Therefore $q(h, L(h)) \geq (1-2t)^2 = 1-4u$. Outside H the coefficient is at most $4u$. The full correct-graph score is $F^2 + A \leq F^2 + t^2 = 1-2u$. Removing a direction loses at least $1-4u$; replacing its label by a wrong one loses at least $1-8u$. This proves (II.108.1).

Since $t \leq 1/10$, all these coefficients are positive. Dropping the last two terms gives the infidelity bound. Also

$$z + \beta \leq \frac{\varepsilon - 2u}{1 - 8u} \leq \varepsilon,$$

the last step using $\varepsilon \leq 1/4$. Two distinct homomorphisms agree on at most half of G ; two candidates satisfying the stated error bound would agree on more than half, proving uniqueness.

For phase sharpness take $G = \mathbb{Z}/2\mathbb{Z}$, a nontrivial character χ , $v = \sqrt{1-t} + \sqrt{t}\chi$, and the correct trivial derivative graph. Its score is $(1-t)^2 + t^2$, and the best quadratic fidelity is $1-t$ for these small t . For graph sharpness take $v = 1$, and change the trivial label to a nontrivial one on any prescribed set of directions, or omit those directions. The score is exactly one minus the error fraction. Rational error fractions are attained on finite groups, which also establishes uniform sharpness by approximation. $\square$

[write] **6 October 2026; sources 28–33.** *Against sources 11 and 29.* Source 11's Corollary II.32.2 treats selected graphs over an odd prime N with $|F| \leq 1$, with error $D_0(1-\alpha^4)N$; the theorem here holds on every finite abelian group, with partial graphs and no amplitude bound. Source 29's Theorem II.110.1 has 20ε and 22ε where this theorem has ε and $\varepsilon/2 + O(\varepsilon^2)$ (note there).

The corresponding best squared-distance bound is $2(1 - \sqrt{1 - \mathcal{G}(\varepsilon)})$. Here $\mathcal{G}(\varepsilon) = \varepsilon/2 + \varepsilon^2/4 + \varepsilon^3/4 + O(\varepsilon^4)$. No claim of joint optimality is made for the coefficient $1 - 8t(1-t)$ in (II.108.1).

The theorem uses no bound on $\|f\|_\infty$ and no restriction on torsion. On groups with two-torsion the phase must be interpreted in the finite-difference sense: for example, the phase taking values $1, i$ on $\mathbb{Z}/2\mathbb{Z}$ is quadratic. Restricting quadratic phases to expressions that divide by two inside the group would incorrectly exclude such examples.

II.109 The correct stability statement for the weights

Source 31, delivered Section 9.

Near equality of a sharp inequality often suggests closeness to its extremizers. Here the topology is decisive. Even with the input function exactly extremal, the weight need not approach the extremal ray in normalized L^1 .

Fix $\psi \in \mathcal{Q}_{\cos}(G)$ and its maximal isotropic subgroup H . Thus $|H| = N$ and $q_\psi = \mathbf{1}_H$. For $w : V \rightarrow \mathbb{R}_{\geq 0}$ put

$$v_H = w|_H, \quad a = \mathbb{E}_{h \in H} v_H(h), \quad w_{\text{out}} = w \mathbf{1}_{V \setminus H}, \quad D = \mathcal{E}(w) - a^4.$$

For a nonzero coset $z \in V/H$, let $a_z = N^{-1} \sum_{u \in z} w(u)$. The notation v_H here denotes a weight, not the normalized input function $v = f/\sqrt{m}$ used earlier.

Theorem II.109.1 (Uniformity and coset-leakage defect). *For every nonnegative w ,*

$$D \geq \|v_H - a\|_{U^2(H)}^4 + \mathcal{E}(w_{\text{out}}) + 4a^2 \sum_{z \neq 0} a_z^2. \quad (\text{II.109.1})$$

If $a > 0$, equality $D = 0$ holds exactly for $w = a\mathbf{1}_H$. If $a > 0$, $v_H - a$ has at most R nonzero Fourier coefficients, and w occupies at most J nonzero cosets, then

$$\frac{1}{N} \sum_{u \in V} |w(u) - a\mathbf{1}_H(u)| \leq R^{1/4} D^{1/4} + \frac{\sqrt{J}}{2a} D^{1/2}. \quad (\text{II.109.2})$$

Proof. The additive quadruples wholly in H contribute

$$\mathcal{E}(w\mathbf{1}_H) = \sum_{\gamma \in \widehat{H}} |\widehat{v}_H(\gamma)|^4 = a^4 + \|v_H - a\|_{U^2(H)}^4.$$

Those wholly outside contribute $\mathcal{E}(w_{\text{out}})$. Among the remaining nonnegative terms, retain quadruples with exactly two outside positions, one on each side of the additive equation. There are four choices of these positions. The two outside elements must be in the same nonzero coset z . Writing the weights in this coset as a function w_z on H , each choice contributes

$$\sum_{\gamma} |\widehat{v}_H(\gamma)|^2 |\widehat{w}_z(\gamma)|^2 \geq a^2 a_z^2.$$

This proves (II.109.1). If $D = 0$ and $a > 0$, every outside mass vanishes and every nontrivial Fourier coefficient of v_H vanishes, giving the equality classification.

For (II.109.2), Fourier Cauchy–Schwarz and Parseval give

$$\|v_H - a\|_1 \leq \|v_H - a\|_2 \leq R^{1/4} \|v_H - a\|_{U^2(H)}.$$

For the outside term, Cauchy–Schwarz on the J coset masses and (II.109.1) gives $\sum_{z \neq 0} a_z \leq \sqrt{J} D^{1/2} / (2a)$. $\square$

II.109.1 Oscillating weights on the exact graph

Proposition II.109.2 (No uniform L^1 modulus for graph weights). *There is no function $\omega(s) \rightarrow 0$ as $s \downarrow 0$ such that*

$$\mathcal{E}(w) \leq (1+s)a^4 \implies \frac{1}{aN} \sum_{u \in V} |w(u) - a\mathbf{1}_H(u)| \leq \omega(s)$$

for all finite abelian groups, even when w is an indicator supported on H and the input function is the constant phase.

Proof. Let $G = \mathbb{F}_2^n \times \mathbb{F}_2^n$, $N = 2^{2n}$, and $b(x, y) = (-1)^{x \cdot y}$. Put $\alpha = N^{-1/2}$, $H = G \times \{1\}$, and $w(h, \chi) = v_0(h)\mathbf{1}_{\chi=1}$ with $v_0 = (1+b)/2$. Summing first over y shows

$$\widehat{b}(u, v) = \alpha(-1)^{u \cdot v}, \quad \mathbb{E}b = \alpha.$$

Thus $a = (1 + \alpha)/2$, and exact Fourier calculation gives

$$\frac{\mathcal{E}(w)}{a^4} = 1 + \frac{N-1}{N^2(1+\alpha)^4} \longrightarrow 1, \quad \frac{\mathbb{E}|v_0 - a|}{a} = 1 - \alpha \longrightarrow 1. \quad (\text{II.109.3})$$

The displayed implication would contradict these limits. $\square$

This example spreads its oscillation over many small Fourier coefficients. Its U^2 discrepancy tends to zero, exactly as Theorem II.109.1 predicts, while its L^1 discrepancy does not. The complexity factor $R^{1/4}$ cannot simply be omitted.

II.109.2 Diffuse weight outside the extremal subgroup

Fix $L > 0$. Let w be constant on the N cosets of H , with value one on H and value $L/(N-1)$ on every other coset. Then $a = 1$, and the outside mass divided by N is exactly L . Fourier orthogonality on V/H gives

$$\mathcal{E}(w) = \frac{(1+L)^4 + (N-1)(1-L/(N-1))^4}{N} \longrightarrow 1. \quad (\text{II.109.4})$$

For all sufficiently large N , these weights lie in $[0, 1]$. Even $L = N^c$, $0 < c < 1/4$, makes the defect tend to zero while the normalized outside mass tends to infinity. This independently shows why an ℓ^2 coset profile does not control its ℓ^1 mass without a bound on the number of occupied cosets.

II.109.3 Transfer to a nearby input

Lemma II.109.3 (Bounded-weight measurement transfer). *For unit vectors v, ψ and $0 \leq w \leq 1$,*

$$|\Theta_w(v) - \Theta_w(\psi)| \leq \sqrt{1 - |\langle v, \psi \rangle|^4}. \quad (\text{II.109.5})$$

Proof. The ambiguity arrays A_v, A_ψ have norm one with respect to normalized h and counting χ measure, and

$$\frac{1}{N} \sum_{h, \chi} A_v(h, \chi) \overline{A_\psi(h, \chi)} = |\langle v, \psi \rangle|^2.$$

Both assertions follow from Parseval and the bijection $(h, x) \mapsto (x, x - h)$. The arrays of squared moduli, divided by N , are probabilities. For any probabilities p, q ,

$$\frac{1}{2} \sum |p - q| \leq \sqrt{1 - \left(\sum \sqrt{pq} \right)^2},$$

by Cauchy-Schwarz applied to $|\sqrt{p} - \sqrt{q}|(\sqrt{p} + \sqrt{q})$. The triangle inequality bounds the sum of square roots below by the modulus of the ambiguity-array inner product. Finally a weight in $[0, 1]$ changes its expectation by at most total variation, proving the lemma. $\square$

Proposition II.109.4 (Exact transfer to a coset phase). *If $\psi \in \mathcal{Q}_{\text{cos}}(G)$, let H be its extremal subgroup and let p_γ be the coefficient probabilities of a unit vector v in its common eigenbasis. Label ψ by zero, put $F = p_0 = |\langle v, \psi \rangle|^2$, and set $A = \sum_{\gamma \neq 0} p_\gamma^2$. Then*

$$\sup_{0 \leq w \leq 1} |\Theta_w(v) - \Theta_w(\psi)| = 1 - \sum_{\gamma} p_\gamma^2 = 1 - F^2 - A \leq 1 - F^2. \quad (\text{II.109.6})$$

There is no small-error hypothesis in this identity.

Proof. The supremum is the total variation distance between the probabilities q_v/N and q_ψ/N . On H one has $q_v \leq 1 = q_\psi$, while outside H one has $q_\psi = 0$. The mass identity therefore gives

$$\frac{1}{2N} \sum_{z \in V} |q_v(z) - \mathbf{1}_H(z)| = \frac{1}{N} \sum_{z \notin H} q_v(z) = 1 - \frac{1}{N} \sum_{z \in H} q_v(z) = 1 - \sum_{\gamma} p_{\gamma}^2.$$

Taking $w = \mathbf{1}_H$ attains this absolute difference. □

For example, suppose $0 \leq w \leq 1$, $\delta < 1/120$, $0 \leq \eta < 1$, and

$$\Theta_w(v)^4 \geq (1 - \eta)\kappa(v)\mathcal{E}(w).$$

Choose the coset phase supplied by Theorem II.107.4. Writing $F = |\langle v, \psi \rangle|^2$, $a = \Theta_w(\psi)$ and $\eta_* = [(1 - \eta)\kappa]^{-1} - 1$, we obtain

$$0 \leq \mathcal{E}(w) - a^4 \leq \eta_* + 4(1 - F^2 - A) \leq \eta_* + 4(1 - F^2). \quad (\text{II.109.7})$$

Indeed both scores lie in $[0, 1]$, so their fourth powers differ by at most four times their difference. Theorem II.109.1 now supplies a coupled stability statement in the topology it actually controls. Since $1 - F \leq \mathcal{R}(\delta)$, a convenient bound depending only on the two input errors is

$$\mathcal{E}(w) - a^4 \leq \frac{1}{(1 - \eta)(1 - \delta)} - 1 + 8\mathcal{R}(\delta) - 4\mathcal{R}(\delta)^2.$$

This is $\eta + 3\delta + O(\eta^2 + \eta\delta + \delta^2)$ as both errors tend to zero. The exact transfer identity is sharp; the subsequent bound for the coupled energy defect is not claimed optimal.

II.110 A joint endpoint rigidity theorem

Source 29, delivered Section 5.

Let G be a finite abelian group, with normalized averaging on G and counting measure on its character group $\hat{G}$. Put

$$A_f(h, \chi) = \mathbb{E}_x f(x) \overline{f(x - h)\chi(x)}, \quad T_\phi(f) = \mathbb{E}_h |A_f(h, \phi(h))|^2,$$

where $\phi : G \rightarrow \hat{G}$ is arbitrary. A quadratic phase means a map $Q : G \rightarrow \mathbb{S}^1$ all of whose third multiplicative derivatives are identically one. Its derivative frequency map is the unique map $\psi : G \rightarrow \hat{G}$ for which $\Delta_h Q = c_h \psi(h)$ with $|c_h| = 1$. This map exists: every first derivative of a quadratic phase has all second derivatives equal to one, and every such unit phase is a constant multiple of a character. The latter assertion follows by evaluating its constant first derivatives at zero. Uniqueness follows from character orthogonality. The derivative cocycle then shows that ψ is a homomorphism.

Theorem II.110.1 (Explicit joint endpoint rigidity). *Suppose $f : G \rightarrow \mathbb{C}$ has $m = \mathbb{E}|f|^2 > 0$, and*

$$T_\phi(f) \geq (1 - \varepsilon)m^2, \quad 0 \leq \varepsilon \leq 2^{-9}.$$

There are a quadratic phase Q , its homomorphic derivative frequency map $\psi : G \rightarrow \widehat{G}$, and $c \in \mathbb{C}$ with $|c|^2 = m$ such that

$$\mathbb{P}_h[\phi(h) \neq \psi(h)] \leq 20\varepsilon, \quad \mathbb{E}_x|f(x) - cQ(x)|^2 \leq 22\varepsilon m.$$

The proof also gives

$$T_\psi(f) \geq (1 - 21\varepsilon)m^2.$$

Thus each displayed conclusion remains valid with the simpler common constant 24 where applicable. No restriction on the exponent, rank, or order of G is imposed, and there is no pointwise bound on f .

[write] **6 October 2026; sources 28–33.** *Constants superseded the same day by source 31* (checked at the write). With $v = f/\sqrt{m}$ the normalized score of source 31’s Section II.108 is $T_\phi(f)/m^2$, so its Theorem II.108.1 applies under the hypothesis here on the larger range $\varepsilon \leq 1/400$: the graph error is at most ε (here 20ε), and the quadratic infidelity at most $(1 - \sqrt{1 - 2\varepsilon})/2 = \varepsilon/2 + O(\varepsilon^2)$, so the squared distance $\mathbb{E}|f - cQ|^2/m$ is at most $2(1 - \sqrt{1 - \varepsilon/2 + O(\varepsilon^2)}) = \varepsilon/2 + O(\varepsilon^2)$ (here 22ε); both bounds of source 31 are separately sharp. The soft-selector, mixed and distance-equivalence results of Sections II.111–II.112 have no counterpart in source 31. Written from the same pin; neither source cites the other.

[write] **6 October 2026; sources 34–39.** The symmetry step of this section (“Forcing symmetry without dividing by two”, gap 1/2) has a sharp-in- p vector-space analogue in source 36’s Theorem F.41.1 at $k = 1$; source 36 does not cite this source.

II.110.1 Amplitude control and projection

By scaling it suffices to prove the theorem when $m = 1$. Parseval gives

$$\mathbb{E}_h \sum_\chi |A_f(h, \chi)|^2 = \mathbb{E}_{h,x} |f(x)|^2 |f(x - h)|^2 = 1.$$

Also, for every (h, χ) , Cauchy–Schwarz gives $|A_f(h, \chi)| \leq 1$. Put $a = |f|$. The triangle inequality, followed by the Fourier identity for autocorrelations, implies

$$\begin{aligned} T_\phi(f) &\leq \mathbb{E}_h (\mathbb{E}_x a(x) a(x - h))^2 \\ &= \sum_\chi |\widehat{a}(\chi)|^4 \leq (\mathbb{E} a)^2 \sum_\chi |\widehat{a}(\chi)|^2 = (\mathbb{E} a)^2. \end{aligned}$$

The penultimate step uses $a \geq 0$, whence $|\widehat{a}(\chi)| \leq \widehat{a}(1) = \mathbb{E} a$. Consequently $(\mathbb{E} a)^2 \geq 1 - \varepsilon$.

Define $q(x) = f(x)/a(x)$ when $a(x) > 0$, and choose $q(x) = 1$ when $a(x) = 0$. Thus $|q| = 1$ and $f = aq$ everywhere. The exact derivative distance is

$$\begin{aligned} \mathbb{E}_{h,x} |\Delta_h f(x) - \Delta_h q(x)|^2 &= \mathbb{E}_{x,y} (a(x)a(y) - 1)^2 \\ &= 2 - 2(\mathbb{E} a)^2 \leq 2\varepsilon. \end{aligned}$$

In the Hilbert space of functions on $G_h \times G_x$, let P_ϕ be the orthogonal projection which, separately in each h -fiber, retains only the character $\phi(h)$. Then

$$\|(I - P_\phi)\Delta f\|_2^2 = 1 - T_\phi(f) \leq \varepsilon.$$

Both Δf and Δq have squared Hilbert norm one. Hence

$$\begin{aligned}\eta &:= 1 - T_\phi(q) = \|(I - P_\phi)\Delta q\|_2^2 \\ &\leq (\sqrt{\varepsilon} + \sqrt{2\varepsilon})^2 = (3 + 2\sqrt{2})\varepsilon < 6\varepsilon\end{aligned}$$

unless $\varepsilon = 0$, in which case $\eta = 0$. This use of the complementary projection is what avoids an unnecessary square-root loss in the final stability rate.

II.110.2 The derivative cocycle gives a homomorphism test

For each h put $a_h = A_q(h, \phi(h))$, and choose a unit scalar c_h such that $c_h = a_h/|a_h|$ when $a_h \neq 0$ (any unit scalar otherwise). Set $v_h(x) = c_h\phi(h)(x)$ and

$$d_h = \|\Delta_h q - v_h\|_2, \quad d_h^2 = 2(1 - |a_h|) \leq 2(1 - |a_h|^2).$$

Thus $\mathbb{E}_h d_h^2 \leq 2\eta$. The exact cocycle identity is

$$\Delta_{h+k} q(x) = \Delta_h q(x) \Delta_k q(x - h).$$

All factors here and in v_h are unimodular. The triangle inequality therefore gives

$$\|v_{h+k} - v_h v_k(\cdot - h)\|_2 \leq d_{h+k} + d_h + d_k.$$

If $\phi(h+k) \neq \phi(h)\phi(k)$, the two unit character multiples on the left have different frequencies and are orthogonal, so their distance is exactly $\sqrt{2}$. It follows that

$$\mathbf{1}_{\phi(h+k) \neq \phi(h)\phi(k)} \leq \frac{1}{2}(d_{h+k} + d_h + d_k)^2 \leq \frac{3}{2}(d_{h+k}^2 + d_h^2 + d_k^2).$$

Averaging proves

$$\delta := \mathbb{P}_{h,k}[\phi(h+k) \neq \phi(h)\phi(k)] \leq 9\eta \leq 54\varepsilon.$$

II.110.3 A complete elementary repair lemma

Lemma II.110.2 (Majority repair). *Let G, H be finite abelian groups, written additively, and let $\phi : G \rightarrow H$ satisfy*

$$\mathbb{P}_{x,y}[\phi(x+y) \neq \phi(x) + \phi(y)] = \delta < 1/6.$$

There is a homomorphism $\psi : G \rightarrow H$ such that

$$\mathbb{P}_x[\phi(x) \neq \psi(x)] \leq \frac{\delta}{1 - 2\delta}.$$

Proof. For fixed x define $D_x(y) = \phi(x+y) - \phi(y)$. If both additive tests on the pairs $(x+y, z-y)$ and $(y, z-y)$ hold, then $D_x(y) = D_x(z)$. As (y, z) ranges uniformly over G^2 , each of these two ordered pairs is uniform in G^2 . Thus, for every fixed x ,

$$\mathbb{P}_{y,z}[D_x(y) \neq D_x(z)] \leq 2\delta.$$

Writing $p_b = \mathbb{P}_y[D_x(y) = b]$, one has $\sum_b p_b^2 \geq 1 - 2\delta$ and hence $\max_b p_b \geq 1 - 2\delta > 1/2$. Let $\psi(x)$ be this unique majority value. For all x, y we have the identity

$$D_{x_1+x_2}(y) = D_{x_1}(x_2+y) + D_{x_2}(y).$$

For fixed x_1, x_2 , each of the three terms equals its majority value except on a set of y of measure at most 2δ . Since $6\delta < 1$, some y realizes all three majority values, proving $\psi(x_1 + x_2) = \psi(x_1) + \psi(x_2)$. If $\phi(x) \neq \psi(x)$, the original additive test at that x fails on every y with $D_x(y) = \psi(x)$, a set of measure at least $1 - 2\delta$. Averaging the original test therefore gives $\delta \geq (1 - 2\delta)\mathbb{P}_x[\phi(x) \neq \psi(x)]$. $\square$

Apply this lemma to $H = \widehat{G}$ in multiplicative notation. Since $\varepsilon \leq 1/512$, one has $\delta \leq 54/512 < 1/6$, and

$$\theta := \mathbb{P}_h[\phi(h) \neq \psi(h)] \leq \frac{54\varepsilon}{1 - 108\varepsilon} \leq 69\varepsilon.$$

The crude constant is improved using a standard ingredient of homomorphism testing: every test with exactly one erroneous value must fail. This observation is classical; see the homomorphism-testing work of Blum–Luby–Rubinfeld and Ben-Or–Coppersmith–Luby–Rubinfeld [49, 122]. We include the elementary counting argument to make the constants explicit.

Let $E_x = \mathbf{1}_{\phi(x) \neq \psi(x)}$. For uniformly independent h, k , the three indicators E_h, E_k, E_{h+k} are pairwise independent, because each pair among $h, k, h+k$ is uniformly distributed on G^2 . If exactly one indicator is one, the homomorphism identity for ψ shows that the test for ϕ fails. Inclusion–exclusion therefore gives

$$\begin{aligned} \delta &\geq \mathbb{P}[E_h + E_k + E_{h+k} = 1] \\ &= 3\theta - 6\theta^2 + 3\mathbb{P}[E_h = E_k = E_{h+k} = 1] \geq \theta(3 - 6\theta). \end{aligned}$$

Substituting the successive upper bounds on θ yields

$$\theta \leq \frac{54\varepsilon}{3 - 414\varepsilon} \leq 25\varepsilon, \quad \theta \leq \frac{54\varepsilon}{3 - 150\varepsilon} \leq 20\varepsilon.$$

All denominators are positive in the stated range. The numerical inequalities follow respectively from

$$\frac{54}{1 - 108/512} < 69, \quad \frac{54}{3 - 414/512} < 25, \quad \frac{54}{3 - 150/512} < 20.$$

Since every coefficient of a unit phase has modulus at most one,

$$T_\psi(q) \geq T_\phi(q) - \theta \geq 1 - \eta - \theta \geq 1 - 26\varepsilon.$$

II.110.4 Forcing symmetry without dividing by two

For the repaired homomorphism put $B(h, k) = \psi(h)(k)$; it is a bicharacter. Write $\zeta = 1 - T_\psi(q) \leq 26\varepsilon$. For each h put $u_h(x) = \Delta_h q(x) \overline{\psi(h)(x)}$. This function is unimodular, its mean is $A_q(h, \psi(h))$, and

$$\Delta_k \Delta_h q(x) = B(h, k) \Delta_k u_h(x).$$

Consequently the following identity, with its average over the entire increment variable, is exact:

$$\begin{aligned} \mathbb{E}_{k,x} |\Delta_k \Delta_h q(x) - B(h, k)|^2 &= \mathbb{E}_{k,x} |\Delta_k u_h(x) - 1|^2 \\ &= 2 - 2 \operatorname{Re} \mathbb{E}_{k,x} u_h(x) \overline{u_h(x - k)} \\ &= 2 - 2 |\mathbb{E}_x u_h(x)|^2. \end{aligned}$$

Averaging in h gives 2ζ . Multiplicative derivatives commute, so both $B(h, k)$ and $B(k, h)$ can be compared with the same function $\Delta_k \Delta_h q = \Delta_h \Delta_k q$. The squared triangle inequality therefore gives

$$\begin{aligned} \mathbb{E}_{h,k} |B(h, k) - B(k, h)|^2 &\leq 2\mathbb{E}_{h,k,x} |B(h, k) - \Delta_k \Delta_h q(x)|^2 \\ &\quad + 2\mathbb{E}_{h,k,x} |\Delta_h \Delta_k q(x) - B(k, h)|^2 \\ &\leq 8\zeta \leq 208\varepsilon < 1. \end{aligned}$$

Retaining this averaged identity avoids the loss incurred by estimating each derivative separately in norm. If the bicharacter $D(h, k) = B(h, k) \overline{B(k, h)}$ were nontrivial, its radical

$$R = \{h : D(h, k) = 1 \text{ for all } k\}$$

would be a proper subgroup, so $|R|/|G| \leq 1/2$. Character orthogonality yields

$$\mathbb{E}_{h,k} D(h, k) = |R|/|G|, \quad \mathbb{E}_{h,k} |B(h, k) - B(k, h)|^2 = 2(1 - |R|/|G|) \geq 1,$$

contradiction. Hence B is exactly symmetric.

II.110.5 Explicit integration on every finite abelian group

Write $G = \bigoplus_{i=1}^r (\mathbb{Z}/n_i \mathbb{Z}) e_i$, and put $z_{ij} = B(e_i, e_j)$. Then $z_{ij} = z_{ji}$ and $z_{ij}^{n_i} = z_{ij}^{n_j} = 1$. Choose $u_i \in \mathbb{S}^1$ with

$$u_i^{n_i} = z_{ii}^{-n_i(n_i-1)/2}.$$

For integer coordinates define

$$Q_0 \left(\sum_i x_i e_i \right) = \prod_i u_i^{x_i} z_{ii}^{x_i(x_i-1)/2} \prod_{i < j} z_{ij}^{x_i x_j}.$$

Increasing x_i by n_i changes the right side by

$$u_i^{n_i} z_{ii}^{n_i x_i + n_i(n_i-1)/2} \prod_{j \neq i} z_{ij}^{n_i x_j} = 1,$$

so the definition descends to G . Direct multiplication shows

$$Q_0(x + y) = Q_0(x) Q_0(y) B(x, y).$$

Therefore

$$\Delta_h Q_0(x) = Q_0(h) B(h, h)^{-1} \psi(h)(x).$$

This explicitly includes even-order groups. The roots u_i may have order larger than n_i , which is precisely why a classical polynomial with denominator dividing the group exponent must not be imposed.

II.110.6 The last Fourier concentration step

The earlier graph comparison also applies directly to f , because $|A_f(h, \chi)| \leq 1$. Thus

$$T_\psi(f) \geq T_\phi(f) - \theta \geq 1 - 21\varepsilon.$$

Let $g = f\overline{Q_0}$. The derivative formula for Q_0 gives

$$T_\psi(f) = \mathbb{E}_h \left| \mathbb{E}_x g(x) \overline{g(x-h)} \right|^2 = \sum_\chi |\widehat{g}(\chi)|^4.$$

Parseval gives $\sum_\chi |\widehat{g}(\chi)|^2 = 1$. Setting $p_\chi = |\widehat{g}(\chi)|^2$ therefore gives $\max p_\chi \geq \sum p_\chi^2 \geq 1 - 21\varepsilon$. Choose a maximizing χ_0 and a unit scalar c aligned with $\widehat{g}(\chi_0)$. Set $Q = Q_0\chi_0$. Multiplication by a character does not change the derivative frequency map ψ . Finally,

$$\mathbb{E}|f - cQ|^2 = 2 - 2|\widehat{g}(\chi_0)| \leq 2(1 - \sqrt{1 - 21\varepsilon}) \leq 22\varepsilon.$$

For the last inequality it suffices that $\sqrt{1 - 21\varepsilon} \geq 10/11$, which follows from $\varepsilon \leq 1/512$. Rescaling proves the theorem.

II.111 Soft selectors

Source 29, delivered Section 6.

Suppose $w : G \times \widehat{G} \rightarrow [0, \infty)$ satisfies $\sum_\chi w(h, \chi) \leq 1$ for every h , and put

$$T_w(f) = \mathbb{E}_h \sum_\chi w(h, \chi) |A_f(h, \chi)|^2.$$

Theorem II.111.1 (Joint stability for soft selectors). *Suppose $m = \mathbb{E}|f|^2 > 0$ and*

$$T_w(f) \geq (1 - \varepsilon)m^2, \quad 0 \leq \varepsilon \leq 2^{-9}.$$

There are a unit quadratic phase Q , its derivative homomorphism $\psi : G \rightarrow \widehat{G}$, and a scalar c with $|c|^2 = m$, such that

$$\mathbb{E}_x |f(x) - cQ(x)|^2 \leq 22\varepsilon m, \quad \mathbb{E}_h [1 - w(h, \psi(h))] \leq 22\varepsilon.$$

In particular, the common constant 24 is valid for both errors.

Proof. Choose $\phi(h)$ maximizing $|A_f(h, \chi)|^2$. Then $T_\phi(f) \geq T_w(f)$. Theorem II.110.1 yields Q, ψ, c with the stated function error and $T_\psi(f) \geq (1 - 21\varepsilon)m^2$. Since $w(h, \chi) \leq 1$,

$$\begin{aligned} T_w(f) &\leq m^2 \mathbb{E}_h w(h, \psi(h)) + \mathbb{E}_h \sum_{\chi \neq \psi(h)} |A_f(h, \chi)|^2 \\ &\leq m^2 \mathbb{E}_h w(h, \psi(h)) + 21\varepsilon m^2. \end{aligned}$$

Consequently

$$\mathbb{E}_h [1 - w(h, \psi(h))] \leq 22\varepsilon, \quad \mathbb{E}_x |f(x) - cQ(x)|^2 \leq 22\varepsilon m.$$

□

The first defect counts both omitted row mass and mass assigned to incorrect frequencies. The substochastic row hypothesis is part of the theorem. A classification of near equality in the unrestricted relative-energy inequality remains a separate question.

II.111.1 Two-sided comparison with distance to the joint model

Corollary II.111.2 (Joint distance equivalence). *Assume $m = 1$. For a unit quadratic phase Q with derivative map ψ_Q define*

$$d_Q = \inf_{|c|=1} \mathbb{E}|f - cQ|^2 = 2 - 2|\langle f, Q \rangle|, \quad s_Q = \mathbb{E}_h[1 - w(h, \psi_Q(h))], \quad D = \inf_Q(d_Q + s_Q).$$

If $\varepsilon = 1 - T_w(f) \leq 2^{-9}$, then

$$\frac{\varepsilon}{2} \leq D \leq 44\varepsilon.$$

The upper constant 48 is a convenient simplification. The lower bound holds for the full range $0 \leq \varepsilon \leq 1$.

Proof. Since $|A_f(h, \chi)|^2 \leq 1$,

$$T_w(f) \geq T_{\psi_Q}(f) - s_Q.$$

Quadratic twisting and the fourth Fourier moment give

$$T_{\psi_Q}(f) \geq |\langle f, Q \rangle|^4 = (1 - d_Q/2)^4 \geq 1 - 2d_Q,$$

where the last inequality is Bernoulli's inequality and uses $0 \leq d_Q \leq 2$. Thus $\varepsilon \leq 2d_Q + s_Q \leq 2(d_Q + s_Q)$. Taking the infimum gives the lower bound. The two error estimates of Theorem II.111.1 give the upper bound. $\square$

II.111.2 An explicit near-maximal U^3 corollary

Corollary II.111.3 (Explicit near-maximal U^3 stability). *If $|f| \leq 1$ and $\|f\|_{U^3}^8 \geq 1 - \varepsilon$, where $0 \leq \varepsilon \leq 2^{-9}$, there are a unit quadratic phase Q and a unit scalar c such that*

$$\mathbb{E}|f - cQ|^2 \leq 24\varepsilon.$$

[write] **6 October 2026; sources 28–33.** Source 31's Theorem II.107.4 gives the sharp profile without the bound $|f| \leq 1$: infidelity at most $(1 - \sqrt{(1 + \sqrt{1 - 4\delta})/2})/2 = \delta/4 + O(\delta^2)$ for $\delta = 1 - \kappa < 1/120$, on every finite abelian group.

Proof. Put $U = \|f\|_{U^3}^8$ and $m = \mathbb{E}|f|^2 \leq 1$, and choose $\phi(h)$ maximizing $|A_f(h, \chi)|^2$. The Fourier mass of each derivative satisfies

$$\sum_{\chi} |A_f(h, \chi)|^2 = \mathbb{E}_x |f(x)|^2 |f(x - h)|^2 \leq m.$$

Consequently

$$U = \mathbb{E}_h \sum_{\chi} |A_f(h, \chi)|^4 \leq m T_{\phi}(f), \quad T_{\phi}(f)/m^2 \geq U/m^3 \geq 1 - \varepsilon.$$

Apply Theorem II.110.1 to $f/\sqrt{m}$. The already proved inequalities $U \leq m T_{\phi}(f) \leq m^3$ imply $m \geq 1 - \varepsilon$, and the selected phase Q has $|\langle f/\sqrt{m}, Q \rangle| \geq \sqrt{1 - 21\varepsilon}$. After rotating Q , its squared distance to f is at most

$$\begin{aligned} m + 1 - 2\sqrt{m}\sqrt{1 - 21\varepsilon} &= (1 - \sqrt{m})^2 + 2\sqrt{m}(1 - \sqrt{1 - 21\varepsilon}) \\ &\leq \varepsilon^2 + 22\varepsilon \leq 23\varepsilon. \end{aligned}$$

This implies the stated bound. $\square$

The qualitative near-extremizer conclusion is established mathematically; the displayed constant and error normalization make it explicit in the present finite-group setting [7].

II.112 Mixed endpoint rigidity and affine derivative selectors

Source 29, delivered Section 7.

For two functions define

$$A_{f,g}(h, \chi) = \mathbb{E}_x f(x) \overline{g(x-h)\chi(x)}, \quad T_w(f, g) = \mathbb{E}_h \sum_{\chi} w(h, \chi) |A_{f,g}(h, \chi)|^2.$$

Suppose $m_f = \mathbb{E}|f|^2 > 0$, $m_g = \mathbb{E}|g|^2 > 0$ and the nonnegative selector satisfies $\sum_{\chi} w(h, \chi) \leq 1$. Then $T_w(f, g) \leq m_f m_g$ by global Parseval. A mixed endpoint generalization identifies a common quadratic part and a possible relative linear twist.

Theorem II.112.1 (Mixed joint endpoint). *If*

$$T_w(f, g) \geq (1 - \varepsilon) m_f m_g, \quad 0 \leq \varepsilon \leq 2^{-12},$$

there are a unit quadratic phase Q , its derivative homomorphism $\psi : G \rightarrow \widehat{G}$, a character $\rho \in \widehat{G}$, and scalars c_f, c_g with $|c_f|^2 = m_f$, $|c_g|^2 = m_g$, such that

$$\begin{aligned} \mathbb{E}|f - c_f Q|^2 &\leq 192\varepsilon m_f, \\ \mathbb{E}|g - c_g Q \bar{\rho}|^2 &\leq 192\varepsilon m_g, \\ \mathbb{E}_h[1 - w(h, \rho\psi(h))] &\leq 192\varepsilon. \end{aligned}$$

In particular, the recovered mixed selector is affine rather than necessarily a homomorphism. The theorem is valid on every finite abelian group without pointwise bounds on either function.

Proof. Normalize $m_f = m_g = 1$. Since the total selector mass $\mathbb{E}_h \sum_{\chi} w(h, \chi)$ is at most one, some pair (h_0, χ_0) satisfies $|A_{f,g}(h_0, \chi_0)|^2 \geq 1 - \varepsilon$. Choose a unit scalar c aligned with this coefficient. With $(T_h u)(x) = u(x - h)$ and $(M_{\chi} u)(x) = \chi(x)u(x)$, put

$$g_0 = \bar{c} T_{-h_0} M_{\chi_0}^{-1} f, \quad g_0(y) = \bar{c} \overline{\chi_0(y + h_0)} f(y + h_0).$$

The inverse modulation and translation are unitary, so

$$d_0 := \|g - g_0\|_2^2 = \|f - c M_{\chi_0} T_{h_0} g\|_2^2 = 2 - 2|A_{f,g}(h_0, \chi_0)| \leq 2\varepsilon.$$

Both mixed coefficient arrays have total squared norm one in the Hilbert space with averaging in h and counting measure in χ . Parseval also gives their exact squared distance

$$\begin{aligned} \mathbb{E}_h \sum_{\chi} |A_{f,g}(h, \chi) - A_{f,g_0}(h, \chi)|^2 &= \mathbb{E}_{h,x} |f(x)|^2 |g(x-h) - g_0(x-h)|^2 \\ &= d_0. \end{aligned}$$

Since each $w(h, \chi) \leq 1$, multiplication by $\sqrt{1 - w(h, \chi)}$ is a contraction. Applying its triangle inequality to the two coefficient arrays yields

$$1 - T_w(f, g_0) \leq (\sqrt{1 - T_w(f, g)} + \sqrt{d_0})^2 \leq 6\varepsilon.$$

This is again a complementary-projection argument, with a positive contraction replacing an orthogonal projection.

Set

$$w_0(k, \xi) = w(k + h_0, \xi \chi_0).$$

The rows of w_0 still have sum at most one. Direct substitution gives

$$A_{f,g_0}(h, \chi) = c \chi_0(h_0 - h) A_f(h - h_0, \chi \chi_0^{-1}),$$

and hence $T_{w_0}(f) = T_w(f, g_0) \geq 1 - 6\varepsilon$. Because $6\varepsilon \leq 6 \cdot 2^{-12} < 2^{-9}$, the soft self-endpoint theorem applies with defect parameter 6ε . Using its common constant 24, there are Q, ψ, c_f satisfying

$$\|f - c_f Q\|_2^2 \leq 144\varepsilon, \quad \mathbb{E}_k[1 - w_0(k, \psi(k))] \leq 144\varepsilon.$$

Define $\rho = \chi_0 \psi(h_0)^{-1}$. Replacing k by $h - h_0$ gives

$$\psi(h - h_0) \chi_0 = \rho \psi(h),$$

so the last estimate is the claimed affine-selector conclusion, with the stronger constant 144.

It remains to identify the second phase. The unitary transform defining g_0 carries $c_f Q$ to

$$\bar{c} c_f \overline{\chi_0(x + h_0)} Q(x + h_0).$$

The derivative identity for Q shows that $Q(x + h_0)/Q(x)$ is a unit constant times $\psi(h_0)(x)$. Therefore the displayed phase is $c_g Q(x) \overline{\rho(x)}$ for a unit scalar c_g . The triangle inequality now gives

$$\|g - c_g Q \bar{\rho}\|_2^2 \leq (\sqrt{d_0} + \|f - c_f Q\|_2)^2 \leq (\sqrt{2} + \sqrt{144})^2 \varepsilon < 180\varepsilon < 192\varepsilon.$$

The first function error and selector error already have constant 144, so 192 works for all three. Rescaling proves the stated form. $\square$

At $\varepsilon = 0$, this classifies exact mixed absolute endpoints: both functions have a common quadratic phase, their ratio is a scalar times a character, and every selector row has total mass one on the corresponding affine derivative graph. This is an endpoint of the substochastic selector problem, not a classification of equality in an arbitrarily normalized weighted-energy inequality.

II.113 Sharpness of the error order

Source 29, delivered Section 8.

On $G = \mathbb{Z}/2\mathbb{Z}$, let χ be its nontrivial character and take $0 < t < 1/2$,

$$f = \sqrt{1-t} + \sqrt{t} \chi, \quad \psi(h) = 1.$$

Then $\mathbb{E}|f|^2 = 1$ and

$$T_\psi(f) = (1-t)^2 + t^2 = 1 - \varepsilon, \quad \varepsilon = 2t(1-t).$$

The function f is strictly positive. Therefore its best squared distance to any unit phase is attained by the constant phase one and is

$$\inf_{|Q|=1} \mathbb{E}|f - Q|^2 = 2 - 2\mathbb{E}f = 2 - 2\sqrt{1-t} \sim \varepsilon/2.$$

No $o(\varepsilon)$ replacement of the squared L^2 error is possible. For the selector error, take $f = 1$, a nontrivial character χ_0 , and

$$w(h, \chi) = (1 - \varepsilon)\mathbf{1}_{\chi=1} + \varepsilon\mathbf{1}_{\chi=\chi_0}.$$

Then $T_w(f) = 1 - \varepsilon$ and the correct-graph mass defect is exactly ε . These show optimality of the linear order, not of the numerical constant 24.

Relation to established extremizer theory. Polynomial-phase near-extremizers of Gowers norms are studied by Eisner–Tao [7]. The mixed array $A_{f,g}$ is a finite short-time Fourier transform, whose exact minimum-support geometry is classified by Nicola [123]. The statements here quantify joint function and selector recovery under the specified substochastic row hypothesis; the exact extremizer geometries are not asserted as new.

II.114 Source 31’s chapter: weights in the Section 7 restriction [write]

Added in the sixth write, 6 October 2026.

The next four sections are source 31’s Sections 6–8 and Appendix A: an elementary weighted energy extraction, a weighted Freiman restriction through the actual vertical defects, the composition that keeps the selected derivative correlation through an exact restriction, and a weighted square-root mass theorem from Reiher–Schoen. *Logically the chapter follows source 10’s Sections II.20–II.21*, which are its spine, and source 24’s chapter (Section II.72); it is printed here because no number printed earlier may move. Source 31 could not see source 10’s text and credits Green and Tao and source 03 for the mechanism.

What is proved, and what is new in this report. Second routes: the restriction of Theorem II.116.1 is source 10’s Theorem II.21.1 with weights (source 10’s progression is never shorter, proved in the note there); the order-eight constant $2^{-363}c^{171}$ (Theorem II.116.2) is weaker than source 10’s $2^{-209}\gamma^{103}$, and the Reiher–Schoen bound is source 10’s alternative. New here: arbitrary nonnegative weights with a fibre cap K throughout; the correlation-preserving exact restriction on $\mathbb{F}_p$ (Theorem II.117.1), which answers source 02’s Subsection 11.1 in a prime-cyclic, bounded-fibre form (item II.1); and the weighted square-root mass theorem (Theorem II.118.1). Compatibility across slices is open.

Relation to the formal statements (checked at HEAD 5d4718772). `lemma_7_5` and `corollary_7_6` (Sections06_07.lean, lines 183, 197) have exact companions (`Proofs07FreimanRestriction.lean`, line 267; `Proofs07AdditiveRestriction.lean`, line 100). The theorems here imply the formal conclusion of `lemma_7_5` with a better constant, as source 10’s do; none is formalized.

II.115 An elementary weighted energy extraction

Source 31, delivered Section 6.

The high-correlation theorem above recovers essentially all structure when the score is close to one. We now turn to arbitrary positive energy. Let A be an abelian group, possibly infinite, and let $w : A \rightarrow [0, 1]$ have finite nonempty support. In this section all sums and energies are unnormalized. Put

$$M = \sum_x w(x), \quad r_w(d) = \sum_x w(x)w(x-d), \quad E(w) = \sum_d r_w(d)^2.$$

Theorem II.115.1 (Elementary weighted BSG). *If $E(w) \geq cM^3$ with $0 < c \leq 1$, put $\rho = c/(2 - c)$. There is $S \subseteq \text{supp } w$ such that*

$$w(S) \geq \frac{\rho M}{2\sqrt{2}}, \quad (\text{II.115.1})$$

$$|S - S| \leq 6912\sqrt{2}\rho^{-5}c^{-4}M. \quad (\text{II.115.2})$$

In particular $w(S) \geq cM/8$ and $|S - S| \leq 2^{19}c^{-9}M$.

Proof. Let $D = \{d : r_w(d) \geq cM/2\}$ and $J = \sum_{d \in D} r_w(d)$. Since $r_w(d) \leq M$ and $\sum_d r_w(d) = M^2$,

$$|D| \leq 2M/c, \quad cM^3 \leq MJ + (cM/2)(M^2 - J),$$

so $J \geq \rho M^2$. Give $X = \text{supp } w$ the probability measure $\mu(x) = w(x)/M$, and join x, y when $x - y \in D$. This is a symmetric relation of weighted edge density at least ρ . Write $N(x)$ for its neighborhood and $t(x, z) = \mu(N(x) \cap N(z))$.

Call a pair bad when $t(x, z) < \rho^2/12$. Choose y with law μ and put $U = N(y)$, $u = \mu(U)$, and

$$b = (\mu \times \mu)\{(x, z) \in U^2 : (x, z) \text{ is bad}\}.$$

Double counting and Cauchy–Schwarz give $\mathbb{E}u^2 \geq \rho^2$, whereas $\mathbb{E}b \leq \rho^2/12$. Thus some realization has $u^2 - 6b \geq \rho^2/2$. For this realization, $u \geq \rho/\sqrt{2}$ and $b \leq u^2/6$. Delete from U the points whose bad partners in U have μ -measure greater than $u/3$. The surviving set S has measure at least $u - 3b/u \geq u/2$, proving (II.115.1).

For any $a, b \in S$, the common set of good partners has measure at least $u/3$. For each such partner z , the choices $x \in N(a) \cap N(z)$ and $y \in N(b) \cap N(z)$ each have measure at least $\rho^2/12$. Consequently the triples (z, x, y) have total raw weight at least

$$\frac{u\rho^4}{432}M^3 \geq \frac{\rho^5}{432\sqrt{2}}M^3.$$

Each raw weight is at most one. The number of these triples is therefore at least the displayed quantity. The map

$$(z, x, y) \mapsto (a - x, z - x, z - y, b - y) \in D^4$$

is injective for fixed a, b , and its labels determine $a - b = d_1 - d_2 + d_3 - d_4$. Choose one pair representing each member of $S - S$; the corresponding label sets are disjoint. Hence

$$|S - S| \frac{\rho^5 M^3}{432\sqrt{2}} \leq |D|^4,$$

which proves (II.115.2). Finally $\rho \geq c/2$ gives the relaxed constants. $\square$

The proof is a weighted version of the elementary dependent-random-choice route in the earlier project report. Its purpose is a self-contained quantitative interface, not an improvement over the strongest general BSG theorem. Loops and repeated vertices are allowed throughout, so there is no hidden lower bound on M or on the ambient group size.

II.116 A weighted finite form of Gowers' Freiman restriction

Source 31, delivered Section 7.

For a map $\phi : B \rightarrow \mathbb{F}_p$, a Freiman homomorphism of order k means that every equality between sums of k elements of B , with repetitions allowed, is preserved by ϕ . The source domain may be any abelian group.

Theorem II.116.1 (Restriction through the actual vertical defects). *Let p be prime, B finite and nonempty, $\phi : B \rightarrow \mathbb{F}_p$, and Γ its graph. For $k \geq 2$ put*

$$V_k = \{y \in \mathbb{F}_p : (0, y) \in k\Gamma - k\Gamma\}, \quad q = |V_k|.$$

For any nonnegative weights $a(x)$ on B , there is a subset B' on which ϕ is a Freiman homomorphism of order k , with

$$\sum_{x \in B'} a(x) \geq \frac{1 + \lfloor p/(kq) \rfloor}{p} \sum_{x \in B} a(x) \geq \frac{1}{kq} \sum_{x \in B} a(x). \quad (\text{II.116.1})$$

If $q = 1$, the whole domain can be retained. If additionally $|\Gamma - \Gamma| \leq C|\Gamma|$, then

$$\sum_{B'} a \geq \frac{1}{kC^{2k+1}} \sum_B a. \quad (\text{II.116.2})$$

Proof. The defect set is symmetric and contains zero. The case $q = 1$ is immediate. Otherwise put $L = \lfloor p/(kq) \rfloor$. Then $kL \leq p/q < p$. Call $d \in \mathbb{F}_p^\times$ bad if $jd \in V_k \setminus \{0\}$ for some integer $1 \leq j \leq kL$. There are at most $kL(q-1)$ bad directions: for each pair (j, y) there is a unique possible d , because j is nonzero modulo p . If $q < p$, this count is at most $p(q-1)/q < p-1$. If $q = p$, then $L = 0$ since $k \geq 2$. In both cases there is a good nonzero direction d .

The progression $P = \{0, d, \dots, Ld\}$ has $L+1$ distinct points. Average the weight of $\phi^{-1}(t+P)$ over $t \in \mathbb{F}_p$ and choose a translate attaining at least the mean. For two k -tuples from this inverse image having the same domain sum, the difference of their image sums belongs to V_k and equals jd for an integer $|j| \leq kL$. Symmetry and goodness of d force $j = 0$. This proves (II.116.1). The case $p = 2$ is included.

For the small-difference-set consequence, the map

$$V_k \times \Gamma \longrightarrow (k+1)\Gamma - k\Gamma, \quad (y, \gamma) \longmapsto \gamma + (0, y) \quad (\text{II.116.3})$$

is injective. The first coordinate determines γ , because Γ is a graph, and then the second determines y . Plünnecke–Ruzsa gives

$$q|\Gamma| \leq |(k+1)\Gamma - k\Gamma| \leq C^{2k+1}|\Gamma|.$$

Insert this into (II.116.1). □

[write] **6 October 2026; sources 28–33.** *Second route to source 10* (checked at the write). Unweighted, this is source 10's Theorem II.21.1, with the same injection $V_k \times \Gamma \rightarrow (k+1)\Gamma - k\Gamma$ ($V_k = H_k$, $q = h_k$) and the same C^{2k+1} ; source 31 credits Green and Tao and source 03, not source 10, which it could not see (note in Section F.36). Source 10's proof uses the progression $\{0, d, \dots, Ld\}$ with $L = \lfloor (p-2)/(k(h_k-1)) \rfloor$, never shorter than the $\lfloor p/(kq) \rfloor$ here: if $kqL' \leq p$ with $L' \geq 1$, then $k(q-1)L' \leq p - kL' \leq p-2$. (The intake found it strictly longer in 5,001 of 96,663 cases $p < 400$, $k \leq 8$.) New here: arbitrary nonnegative weights.

We use the standard form $|rA - sA| \leq C^{r+s}|A|$ under $|A - A| \leq C|A|$. For completeness, it follows from the minimal-growth lemma of Petridis [36]: choose nonempty $X \subseteq A$ minimizing

$|X - A|/|X| \leq C$; that lemma gives $|X - jA| \leq C^j|X|$. Ruzsa's triangle inequality then gives $|rA - sA| \leq |X - rA||X - sA|/|X| \leq C^{r+s}|A|$. This standard theorem is the only imported additive-growth input in the elementary restriction route.

The exact comparison. Gowers' Lemma 7.5, and its corrected repository statement, retain $|B|/(8kC^{4k})$ in the unweighted case [1, 139]. Theorem II.116.1 increases this particular guarantee by the factor $8C^{2k-1}$ and permits arbitrary nonnegative weights. The smaller sumset in (II.116.3) is not a new mechanism: Green–Tao use the order-eight instance $\Gamma' + V_8 \subseteq 9\Gamma' - 8\Gamma'$ [4]. The exact finite factor in (II.116.1) also records the rounding and includes the singleton-progression regime.

II.116.1 Completing the weighted graph extraction

Theorem II.116.2 (A prescribed weight on an exact Freiman graph). *Let A be an abelian group, let p be prime, and let $w : A \times \mathbb{F}_p \rightarrow [0, 1]$ have finite nonempty support, mass M , and $E(w) \geq cM^3$, where $0 < c \leq 1$. Suppose each first-coordinate fiber of its support contains at most $K \geq 1$ points. For every $k \geq 2$ there is an order- k Freiman map $\phi : B \rightarrow \mathbb{F}_p$ whose graph lies in $\text{supp } w$ and satisfies*

$$\sum_{x \in B} w(x, \phi(x)) \geq \frac{c^{20k+11}}{k 2^{42k+24} K^{2k+2}} M. \quad (\text{II.116.4})$$

For $k = 8$ the fraction is $2^{-363} c^{171} K^{-18}$.

Proof. Use Theorem II.115.1 to obtain S . In each nonempty fiber of S , retain a point of maximum original weight, forming a graph Γ_1 . Its weight and cardinality are at least $\rho M/(2\sqrt{2}K)$. Its difference set is contained in $S - S$, so

$$\frac{|\Gamma_1 - \Gamma_1|}{|\Gamma_1|} \leq 27648K\rho^{-6}c^{-4} \leq 1769472Kc^{-10} < 2^{21}Kc^{-10}.$$

Apply Theorem II.116.1 to Γ_1 with the original weights and use $\rho/(2\sqrt{2}) \geq c/8$. The retained fraction is at least

$$\frac{c}{8K} \cdot \frac{1}{k(2^{21}Kc^{-10})^{2k+1}},$$

which is exactly (II.116.4). $\square$

The argument can retain the actual maximum-weight graph mass and the actual defect count q instead of replacing them by K and C^{2k+1} . The displayed theorem is a uniform statement suitable for a quantitative dependency ledger, not a recommendation to discard those measured improvements.

II.116.2 Order-eight comparison, including the modern input

For an unweighted graph with domain size M and at least cM^3 additive quadruples, $K = 1$. The historical Corollary 7.6 gives $2^{-1882}c^{1164}M$. The elementary route just proved gives $2^{-363}c^{171}M$.

Reiher–Schoen's theorem [12, Theorem 1.2] gives a subset of size at least $(1-\eta)\sqrt{c}M$ and difference ratio at most $2^{33}\eta^{-9}c^{-4}$, for $0 < \eta < 1/2$. Combining this published input with Theorem II.116.1 gives the further valid bound

$$\frac{(1-\eta)\eta^{18k+9}}{k 2^{66k+33}} c^{8k+9/2} M. \quad (\text{II.116.5})$$

For $k = 8$, $\eta = 1/4$, this is $3 \cdot 2^{-872} c^{137/2} M$.

Route	Retained domain	Status
Gowers Corollary 7.6	$2^{-1882} c^{1164} M$	Historical baseline
Elementary weighted route, $K = 1$	$2^{-363} c^{171} M$	Proved here
Reiher–Schoen plus restriction	$3 \cdot 2^{-872} c^{137/2} M$	Derived corollary

The second and third rows are not uniformly ordered because their constants differ. One may take their maximum, or optimize (II.116.5) in η . These are local comparisons; none has been propagated through every stage of Gowers’ proof.

[write] **6 October 2026; sources 28–33.** *Against source 10* (checked at the write). Source 10’s Corollary II.21.2 gives $2^{-(24k+14)} \gamma^{12k+7}/k$, that is $2^{-209} \gamma^{103}$ at order eight, under the same hypothesis ($\gamma = c$, $K = 1$), against $2^{-363} c^{171}$ here; its weighted extraction (Theorem II.20.1) has $(c/4, 2^{10} c^{-5})$ where Theorem II.115.1 has source 03’s $(c/8, 2^{19} c^{-9})$ in effect. The Reiher–Schoen bound (II.116.5) is source 10’s modern-input alternative after its Corollary II.21.2: $\eta^{18k+9} = \eta^{9(2k+1)}$, $2^{66k+33} = 2^{33(2k+1)}$, the same power of c . New here are the weights and the fibre cap K ; source 24’s $2^{-295} \gamma^{103}$ is a third route (Section II.72).

II.117 Keeping the original derivative correlation

Source 31, delivered Section 8.

Now let $G = \mathbb{F}_p$ and identify its characters with $\mathbb{F}_p$. Continue to use the normalized coefficients $q = q_f \in [0, 1]$. Fix any candidate set $\Omega \subseteq \mathbb{F}_p^2$. For $0 < \tau \leq 1$ let

$$\Delta_\tau = \{(h, r) \in \Omega : q(h, r) \geq \tau\}$$

be nonempty, and define its profile

$$\mu = \frac{1}{p} \sum_{\Delta_\tau} q, \quad \nu = \frac{1}{p} \sum_{\Delta_\tau} q^2, \quad b = \max_{\Delta_\tau} q, \quad K = \max_h |\{r : (h, r) \in \Delta_\tau\}|. \quad (\text{II.117.1})$$

These statistics retain more information than a threshold alone. We have $0 < \nu \leq \mu \leq 1$ and $0 < b \leq 1$. For the unrestricted choice $\Omega = \mathbb{F}_p^2$, one has $b = 1$ because $q(0, 0) = 1$; restricting the candidate set allows $b < 1$.

Theorem II.117.1 (Correlation-preserving exact restriction). *Put*

$$c = \frac{\nu^4}{b\kappa\mu^3}. \quad (\text{II.117.2})$$

Then $0 < c \leq 1$. For every $k \geq 2$, there is an order- k Freiman map $\phi : B \rightarrow \mathbb{F}_p$ with graph contained in Δ_τ , for which

$$\frac{1}{p} \sum_{h \in B} q_f(h, \phi(h)) \geq \frac{c^{20k+11}}{k^{242k+24} K^{2k+2}} \mu. \quad (\text{II.117.3})$$

In the original derivative normalization, both sides are multiplied by m^2 . Every selected Fourier coefficient has modulus at least $m\sqrt{\tau}$.

Proof. Take $w = b^{-1}q\mathbf{1}_{\Delta_\tau}$. It lies in $[0, 1]$, has mass $M = p\mu/b$, and score $\Theta_w = \nu/b$. Proposition II.105.2 yields

$$\mathbf{E}(w) \geq \frac{p^3\nu^4}{b^4\kappa} = cM^3.$$

For any weight in $[0, 1]$, one has $r_w(d) \leq M$ and $\sum_d r_w(d) = M^2$, so $\mathbf{E}(w) \leq M^3$. This proves $c \leq 1$; positivity is immediate. Apply Theorem II.116.2, and multiply its weight conclusion by b/p . $\square$

[write] **6 October 2026; sources 28–33.** *Answers source 02's Subsection 11.1 in a prime-cyclic, bounded-fibre form* (item II.1): a fixed fraction, explicit in μ, ν, b, κ, K , of the selected derivative correlation survives to an exact Freiman graph. Compatibility across slices is open (Research question in Section V.123).

For arbitrary f an available fiber cap is

$$K \leq \left\lfloor \frac{R_f}{\tau} \right\rfloor, \quad R_f = \max_h \frac{\mathbb{E}_x |f(x)|^2 |f(x-h)|^2}{m^2}.$$

If $|f| \leq 1$, then $R_f \leq 1/m$. In particular the theorem does not silently assume that every fiber has total normalized mass one; that assertion would fail for nonconstant amplitudes.

The information in κ survives all the way to an exact Freiman map. For instance, with the same μ, ν, b, K , a smaller κ increases the certified relative energy c . This is a valid profile comparison, not a claim that these statistics vary independently for actual functions.

A threshold chosen from the norm. For $\Omega = \mathbb{F}_p^2$ and $\tau = \kappa/2$, the discarded fourth-moment contribution is at most $\tau p^{-1} \sum q = \kappa/2$, so $\nu \geq \kappa/2$. This gives a completely explicit positive-correlation application for every nonzero f . The profile version is preferable because the worst-case powers are large and actual spectra can be much more concentrated. The theorem makes no claim that independently selected graphs from different higher-order slices agree with one another; such compatibility is required for further integration.

II.118 A square-root weighted mass theorem from Reiher–Schoen

Source 31, delivered Appendix A.

This appendix records a stronger weighted mass conclusion available from a published input. It is a derived corollary, not a claim that weighted BSG or probabilistic thickening is new.

Theorem II.118.1 (Weighted square-root mass). *Let A be an abelian group and let $w : A \rightarrow [0, 1]$ have finite nonempty support, mass M , and $\mathbf{E}(w) \geq cM^3$, where $0 < c \leq 1$. For every $0 < \epsilon < 1$ there is $S \subseteq \{x : w(x) \geq \epsilon c/16\}$ such that*

$$w(S) \geq (1 - \epsilon)\sqrt{c}M, \quad |S - S| \leq 2^{60}\epsilon^{-11}c^{-11/2}|S|. \quad (\text{II.118.1})$$

Proof. We first prove a parameterized version. Choose $0 < \eta < 1/2$, $0 < \theta < 1$, and $0 < \tau < c/4$; put $q = c - 4\tau$. Truncate to $v = w\mathbf{1}_{w \geq \tau}$, and write $M_v = \sum v$, $e = \mathbf{E}(v)$, and $n = |\text{supp } v|$. Telescoping the four factors of the energy gives

$$e \geq qM^3, \quad n \leq M/\tau, \quad M_v \leq M.$$

Indeed, in a term containing one discarded factor, bound that factor by τ and eliminate its coordinate from the additive equation. The other three sums are at most M each.

For integers $L \rightarrow \infty$, independently include each point $(x, z) \in \text{supp } v \times \mathbb{Z}/L\mathbb{Z}$ in a random set A_L with probability $v(x)$. There are deterministic choices along which

$$|(A_L)_x|/L \rightarrow v(x), \quad |A_L|/L \rightarrow M_v, \quad \mathbb{E}(A_L)/L^3 \rightarrow e. \quad (\text{II.118.2})$$

Here are explicit second-moment estimates establishing this step. Each fiber size has variance at most L . For every base additive quadruple, all but at most $6L^2$ of its L^3 lifted quadruples have distinct lifted points. Thus

$$|\mathbb{E}\mathbb{E}(A_L) - L^3 e| \leq 6n^3 L^2.$$

Two energy monomials have zero covariance unless they share a lifted point. Choosing the shared positions and the five free points bounds the number of such pairs by $16n^5 L^5$. Therefore $\text{Var}(\mathbb{E}(A_L)) \leq 16n^5 L^5$. Chebyshev and a finite union bound give (II.118.2), for example with normalized deviations $L^{-1/4}$. Only coefficients 1 and -1 were eliminated, so no torsion restriction is used.

Apply Reiher–Schoen to A_L with its actual $K_L = |A_L|^3/\mathbb{E}(A_L)$. It gives $B_L \subseteq A_L$ with

$$|B_L| \geq (1 - \eta)K_L^{-1/2}|A_L|, \quad |B_L - B_L| \leq 2^{33}\eta^{-9}K_L^4|B_L|.$$

Pass to a subsequence so that every fiber ratio $|(B_L)_x|/L$ converges to $b(x)$. Then $0 \leq b \leq v$, and, with $r = \sum b$,

$$r \geq (1 - \eta)\sqrt{e/M_v} \geq (1 - \eta)\sqrt{q} M.$$

The difference-set factors tend to a number $D_0 \leq 2^{33}\eta^{-9}q^{-4}$.

Put $t_0 = \theta(1 - \eta)\tau\sqrt{q}$ and $S = \{x : b(x) \geq t_0\}$. The discarded b -mass is at most $t_0 n \leq \theta r$. Hence $w(S) \geq (1 - \theta)r$ and $|S| \geq (1 - \theta)r$. For every $z \in S - S$, choose a representing pair of fibers. The corresponding fiber of $B_L - B_L$ contains their difference set, of size at least $(t_0 - o(1))L$. Summing over distinct base differences, dividing by L , and taking limits gives $|S - S| \leq D_0 r/t_0$. Consequently

$$\begin{aligned} w(S) &\geq (1 - \eta)(1 - \theta)\sqrt{q} M, \\ \frac{|S - S|}{|S|} &\leq \frac{2^{33}\eta^{-9}}{(1 - \eta)\theta(1 - \theta)\tau q^{9/2}}. \end{aligned}$$

Finally choose $\eta = \theta = \epsilon/4$ and $\tau = \epsilon c/16$. The mass factor is $(1 - \epsilon/4)^{5/2}\sqrt{c} \geq (1 - \epsilon)\sqrt{c}$. The doubling factor is

$$2^{57}\epsilon^{-11}c^{-11/2}(1 - \epsilon/4)^{-13/2} < 2^{60}\epsilon^{-11}c^{-11/2},$$

as required. □

[write] 6 October 2026; sources 28–33. This bears on source 10’s Question 7 (item II.18): it is the weighted form of the Reiher–Schoen retention that question asks to propagate; the propagation itself is not attempted here.

Reiher–Schoen’s indicator examples already show that the square-root order of the retained mass cannot in general be replaced by a larger order with uniform small doubling. The numerical constant and the doubling exponent in (II.118.1) are not claimed optimal. The passage through large auxiliary cyclic groups is an existence proof; it is not presented as an efficient extraction algorithm.

II.119 Source 29 at Lemma 9.3: restriction over a fixed field [write]

Added in the sixth write, 6 October 2026.

The next section is source 29’s Section 9. It continues source 17’s random-linear-map restriction over a fixed characteristic (Section II.63) and source 04’s restriction theorem (Theorem II.9.1), keeping the affine-rank stratification of every tuple, degenerate ones included. Over a fixed field $\mathbb{F}_q$ and a fixed translation-invariant system with s free coordinates the size threshold becomes $8qS/(\alpha\eta)$, linear in $(\alpha\eta)^{-1}$, with retention proportional to $\alpha^{s+1}\eta^s|B|^s$ (Theorem II.120.4), and a weaker coefficient removes the threshold (Corollary II.120.5). This answers source 04’s Question 4 and source 17’s Question 2 for fixed fields only (item II.6). The constants depend on q , so the theorem does not replace the all-moduli `lemma_9_3` (`Sections08_09.lean`, line 77; companion `lemma_9_3_holds`, `Proofs09Restriction.lean`, line 376, at HEAD 5d4718772), as the source says (Section V.115). Not formalized. Source 29’s introduction is Section F.34.

II.120 Restriction without discarding affine degeneracies

Source 29, delivered Section 9.

Attribution and scope. The random graph-linear selector and the exact survival formula for *affine-generic* tuples already occur in Section 8 of *Polynomial-loss restriction in Gowers’s Szemerédi argument* [136]. Its proof charges every nongeneric tuple the upper bound one for its survival probability. The new contribution here is to retain the affine-rank stratification of *all* tuples. This removes any ambient-dimension requirement from the ordinary rank-one version, gives a threshold proportional to $(\alpha\eta)^{-1}$ at fixed field and fixed system, and then removes that threshold entirely at the cost of a smaller constant. This answers the fixed-characteristic version of that manuscript’s proposed rank-stratified threshold problem. It does not solve its separate prime-cyclic Fejér-kernel threshold problem. The constants below depend on the field; their dependence is important when comparing the results.

II.120.1 Setup and the exact survival polynomial

Let $\mathbb{F}_q$ be a finite field, let V, W be finite-dimensional $\mathbb{F}_q$ -vector spaces, and let $\mathcal{R} \leq \mathbb{F}_q^\ell$ be a relation space satisfying

$$\mathcal{R} \subseteq \mathbf{1}^\perp, \quad d = \dim \mathcal{R}, \quad s = \ell - d \geq 2.$$

The first condition says that every prescribed relation has coefficient sum zero. Thus the system is invariant under translating all its entries by the same vector. A tuple $x = (x_1, \dots, x_\ell)$ is a solution when

$$\sum_{i=1}^{\ell} a_i x_i = 0 \quad \text{for every } a \in \mathcal{R}.$$

For a nonempty set $B \subseteq V$ and a map $\phi : B \rightarrow W$, write $T_{\mathcal{R}}(B)$ for the number of ordered solution tuples in B^ℓ and $G_{\mathcal{R}}(B, \phi)$ for the number also satisfying

$$\sum_{i=1}^{\ell} a_i \phi(x_i) = 0 \quad \text{for every } a \in \mathcal{R}.$$

All tuples, including repetitions, are counted. Put $m = |B|$. Row reduction chooses s free coordinate positions, and the other coordinates are uniquely determined by them. Consequently

$$T_{\mathcal{R}}(B) \leq m^s. \quad (\text{II.120.1})$$

The cases $d = 0$ and $s = 1$ are harmless: with no relations every tuple is respected; with $s = 1$, translation invariance forces every solution to be a constant tuple, which every map respects. We keep $s \geq 2$ below.

For each solution tuple define its full affine-relation space and its augmented domain rank by

$$K_x = \left\{ a \in \mathbb{F}_q^\ell : \sum_i a_i = 0, \quad \sum_i a_i x_i = 0 \right\}, \quad r(x) = \ell - \dim K_x.$$

Define the full error map

$$E_x : K_x \longrightarrow W, \quad E_x(a) = \sum_i a_i \phi(x_i), \quad e(x) = \text{rank } E_x.$$

The term “full” matters: a degenerate tuple can respect all the prescribed relations in $\mathcal{R}$ while failing an additional relation in K_x . Thus for degenerate tuples, being respected need not be equivalent to $e(x) = 0$. A tuple is affine-generic exactly when $K_x = \mathcal{R}$, or equivalently $r(x) = s$; on such tuples $e(x)$ is the error rank already used in [136, Section 8].

Choose a uniformly random linear map

$$A : V \oplus W \oplus \mathbb{F}_q \longrightarrow \mathbb{F}_q^t, \quad \lambda = q^{-t},$$

and select

$$B_A = \{x \in B : A(x, \phi(x), 1) = 0\}. \quad (\text{II.120.2})$$

Lemma II.120.1 (Exact survival, including degeneracies). *For every solution tuple x ,*

$$\Pr[x_1, \dots, x_\ell \in B_A] = \lambda^{r(x)+e(x)}.$$

In particular, a respected tuple survives with probability at least λ^s , whereas an unrespected tuple survives with probability at most $\lambda^{r(x)+1}$.

Proof. The relation space among the augmented graph columns $z_i = (x_i, \phi(x_i), 1)$ is precisely $\ker E_x \leq K_x$. Its dimension is $\dim K_x - e(x)$, so

$$\dim \text{span}\{z_1, \dots, z_\ell\} = r(x) + e(x).$$

A uniform linear map to $\mathbb{F}_q^t$ kills a fixed subspace of dimension u with probability $q^{-tu} = \lambda^u$. This event is exactly the simultaneous selection event, whether or not entries repeat. If the tuple is respected, $\mathcal{R} \subseteq \ker E_x$, whence $r(x) + e(x) \leq \ell - \dim \mathcal{R} = s$. If it is unrespected, E_x is nonzero on $\mathcal{R}$, so $e(x) \geq 1$. The final inequalities follow because $0 < \lambda \leq 1$. $\square$

Let $G_{r,e}$ and $U_{r,e}$ count respectively the respected and unrespected solution tuples of full rank pair (r, e) . If $X = G_{\mathcal{R}}(B_A, \phi)$ and $Y = T_{\mathcal{R}}(B_A) - X$, the exact expectation certificate is

$$\mathbb{E}[\eta X - Y] = \sum_{r,e} (\eta G_{r,e} - U_{r,e}) \lambda^{r+e}. \quad (\text{II.120.3})$$

There is no exceptional-set subtraction. Any positive lower bound L for the right side yields a realization with

$$X + Y \geq X \geq L/\eta, \quad \frac{X}{X + Y} \geq \frac{1}{1 + \eta}.$$

The lower bound ensures the denominator is nonzero.

II.120.2 The rank-weighted counting lemma

For $u = (u_1, \dots, u_s) \in B^s$, let

$$r(u) = \dim \text{span}\{(u_i, 1) : 1 \leq i \leq s\}, \quad z = m\lambda, \quad S_{q,s} = \sum_{i=0}^{s-2} q^i = \frac{q^{s-1} - 1}{q - 1}.$$

Lemma II.120.2 (Weighted rank count). *For $0 < \lambda \leq 1$,*

$$R_B(\lambda) := \sum_{u \in B^s} \lambda^{r(u)} \leq z \prod_{i=0}^{s-2} (z + q^i), \quad (\text{II.120.4})$$

$$\sum_{\substack{u \in B^s \\ r(u) < s}} \lambda^{r(u)} \leq z^s \left\{ \prod_{i=0}^{s-2} \left(1 + \frac{q^i}{z} \right) - 1 \right\}. \quad (\text{II.120.5})$$

If $z \geq 2S_{q,s}$, these imply

$$R_B(\lambda) \leq 2\lambda^s m^s, \quad (\text{II.120.6})$$

$$\sum_{r(u) < s} \lambda^{r(u)} \leq 2S_{q,s} \lambda^{s-1} m^{s-1}. \quad (\text{II.120.7})$$

Proof. The first point has m choices and raises augmented rank from zero to one, contributing $m\lambda = z$. Before the j th point is chosen, where $2 \leq j \leq s$, the affine span of the previous points contains at most q^{j-2} points. A choice inside that span does not raise the rank, and has at most q^{j-2} possibilities. A choice outside has at most m possibilities and raises the rank by one, contributing at most $m\lambda = z$. Summing over these two possibilities at each step proves (II.120.4).

More explicitly, classify a tuple by the positions after the first at which the rank rises. The contribution of any particular pattern is at most z times a product, with factor z at a rise and factor q^{j-2} at a non-rise. The tuples of rank less than s have at least one non-rise. Summing only those patterns gives the complete product minus the all-rise term z^s , proving (II.120.5). This subtraction is from the pattern upper bounds; it does not subtract an upper bound for generic tuples from a separate upper bound for all tuples.

For nonnegative a_i with $\sum_i a_i \leq 1/2$, one has

$$\prod_i (1 + a_i) \leq \frac{1}{1 - \sum_i a_i} \leq 2, \quad \prod_i (1 + a_i) - 1 \leq 2 \sum_i a_i.$$

For example, expand the product and bound its degree- j elementary symmetric sum by $(\sum_i a_i)^j$, then sum the geometric series. Use $a_i = q^i/z$ to obtain the two simplified bounds. $\square$

Lemma II.120.3 (Free coordinates preserve affine rank). *Choose the s free positions by row reduction of a basis matrix for $\mathcal{R}$. For every solution tuple, the augmented domain columns are spanned by those in the free positions. Thus the domain rank of the full tuple equals the augmented rank of the free subtuple.*

Proof. Let M_P be the invertible pivot submatrix and M_F the remaining columns. The pivot entries satisfy $x_P = -M_P^{-1}M_F x_F$. Because every row sums to zero, $M_P \mathbf{1}_P + M_F \mathbf{1}_F = 0$. Thus every row of $-M_P^{-1}M_F$ sums to one. The same identities therefore hold for the augmented columns $(x_i, 1)$. The free columns are themselves among all columns, so the two spans are equal. $\square$

II.120.3 A linear threshold in the inverse accuracy

Theorem II.120.4 (Restriction with all degeneracies retained). *Suppose $0 < \alpha, \eta \leq 1$ and*

$$G_{\mathcal{R}}(B, \phi) \geq \alpha m^s.$$

If

$$m \geq \frac{8qS_{q,s}}{\alpha\eta}, \quad (\text{II.120.8})$$

there is a subset $B' \subseteq B$ such that

$$T_{\mathcal{R}}(B') \geq \frac{\alpha^{s+1}\eta^s}{2(4q)^s} m^s, \quad (\text{II.120.9})$$

$$\frac{G_{\mathcal{R}}(B', \phi)}{T_{\mathcal{R}}(B')} \geq \frac{1}{1+\eta} \geq 1-\eta. \quad (\text{II.120.10})$$

One may take B' of the form (II.120.2) with $t = \lceil \log_q(4/(\alpha\eta)) \rceil$.

[write] **6 October 2026; sources 28–33.** *In this report* (checked at the write). This answers source 04's Question 4 (item II.6) and source 17's Question 2 for fixed fields only: over a fixed $\mathbb{F}_q$ the threshold becomes linear in $(\alpha\eta)^{-1}$ (source 04's is $\Lambda_d(\alpha\eta)^{-2d}$, source 17's of degree $2m-1$), and Corollary II.120.5 removes it; source 17's fixed-characteristic theorem (Theorem II.63.5) charges every nongeneric tuple probability one and needs a large domain dimension. As Section V.115 says, the constants depend on q , so this does not replace the all-moduli `lemma_9_3` (Sections08_09.lean, line 77; companion `lemma_9_3_holds`, Proofs09Restriction.lean, line 376, at HEAD 5d4718772); the growing-modulus case stays open (item II.8). Not formalized.

Proof. Put $a = \alpha\eta$, choose the stated t , and set $\lambda = q^{-t}$. Then

$$\frac{a}{4q} < \lambda \leq \frac{a}{4}, \quad m\lambda \geq 2S_{q,s}.$$

Lemma II.120.1 gives $\mathbb{E}X \geq \alpha\lambda^s m^s$. For each unrespected solution tuple, its survival is at most $\lambda^{r(x)+1}$. The free positions determine at most one full solution tuple, and have the same augmented domain rank by Lemma II.120.3. We may sum over all free tuples in B^s , including ones which do not extend to a solution in B^ℓ , to obtain

$$\mathbb{E}Y \leq \lambda R_B(\lambda) \leq 2\lambda^{s+1} m^s.$$

Therefore

$$\mathbb{E}[\eta X - Y] \geq (a - 2\lambda)\lambda^s m^s \geq \frac{a}{2}\lambda^s m^s > 0.$$

Choose a realization at least its expectation. It satisfies $Y \leq \eta X$ and

$$T_{\mathcal{R}}(B_A) \geq X \geq \frac{\alpha}{2}\lambda^s m^s \geq \frac{\alpha^{s+1}\eta^s}{2(4q)^s} m^s.$$

These are the required claims. □

Corollary II.120.5 (No size or dimension threshold). *Under the same input assumption, with any $m > 0$, there is $B' \subseteq B$ such that the respected proportion is at least $1/(1+\eta)$ and*

$$T_{\mathcal{R}}(B') \geq (8qS_{q,s})^{-s} \alpha^{s+1} \eta^s m^s. \quad (\text{II.120.11})$$

Consequently the same subset satisfies

$$|B'| \geq \frac{\alpha^{1+1/s}\eta}{8qS_{q,s}} m. \quad (\text{II.120.12})$$

Proof. If (II.120.8) holds, apply Theorem II.120.4, noting that $(8qS_{q,s})^{-s} \leq [2(4q)^s]^{-1}$ because $s \geq 2$ and $S_{q,s} \geq 1$. Otherwise take any singleton $B' \subseteq B$. Translation invariance makes its unique tuple a respected solution. Moreover

$$(8qS_{q,s})^{-s} \alpha^{s+1} \eta^s m^s = \alpha \left(\frac{\alpha \eta m}{8qS_{q,s}} \right)^s < \alpha \leq 1.$$

This proves the count and ratio in the remaining case. Finally apply (II.120.1) to B' and take the sth root. $\square$

What “no dimension threshold” means. The constants depend on q and s , but not on $\dim V$, $\dim W$, the ambient cardinality, or the density of B in its ambient space. The result remains valid when every tuple is nongeneric. For a growing prime field with $\dim V = 1$, its constants can make it very weak; this does not supersede the bounded-coefficient prime-cyclic construction. The use of vector-space rank is essential to the present proof. No corresponding bound for all composite cyclic moduli, or for arbitrary finite abelian groups of mixed exponent, follows from this argument.

II.120.4 The higher-error-rank benefit survives

Theorem II.120.6 (Generic rank gap with the improved threshold). *In addition to the assumptions of Theorem II.120.4, suppose that every affine-generic unrespected solution tuple has error rank at least h , where $1 \leq h \leq d$. Under the same threshold (II.120.8), there is a subset $B' \subseteq B$ with respected proportion at least $1/(1 + \eta)$ and*

$$T_{\mathcal{R}}(B') \geq \frac{\alpha^{1+s/h} \eta^{s/h}}{2 \cdot 4^{s/h} q^s} m^s. \quad (\text{II.120.13})$$

Proof. Set $a = \alpha \eta$ and choose the least positive integer t such that $\lambda = q^{-t}$ satisfies $\lambda^h \leq a/4$. Then

$$q^{-1}(a/4)^{1/h} < \lambda \leq (a/4)^{1/h}, \quad \lambda > \frac{a}{4q}.$$

The threshold again ensures $m\lambda \geq 2S_{q,s}$. Respected tuples give $\mathbb{E}X \geq \alpha \lambda^s m^s$. Generic unrespected tuples contribute at most $\lambda^{s+h} m^s$ to $\mathbb{E}Y$. Nongeneric unrespected tuples have $r(x) < s$ and $e(x) \geq 1$. Using the free-coordinate injection and (II.120.7), their total expected contribution is at most

$$\lambda \sum_{r(u) < s} \lambda^{r(u)} \leq 2S_{q,s} \lambda^s m^{s-1}.$$

It follows that

$$\mathbb{E}[\eta X - Y] \geq \lambda^s m^s \left(a - \lambda^h - \frac{2S_{q,s}}{m} \right) \geq \frac{a}{2} \lambda^s m^s.$$

Indeed $\lambda^h \leq a/4$ and $2S_{q,s}/m \leq a/(4q) \leq a/4$. The same positive-potential argument as before gives $T_{\mathcal{R}}(B') \geq \alpha \lambda^s m^s / 2$ and the asserted ratio. The lower bound on λ proves (II.120.13). $\square$

The singleton argument does *not* remove the threshold while preserving these stronger exponents for $h > 1$: below the threshold, $\alpha(\alpha\eta)^{s/h} m^s$ need not be bounded by a constant. No such removal is asserted here.

II.120.5 Examples and comparison quantities

For a single balanced additive relation on $2d_0$ variables, with $d_0 \geq 2$, one has $s = 2d_0 - 1$. Write $T_{d_0}(B)$ for the number of ordered tuples satisfying $x_1 + \cdots + x_{d_0} = x_{d_0+1} + \cdots + x_{2d_0}$, and $G_{d_0}(B, \phi)$ for the number also respected by ϕ . Thus the no-threshold count is

$$c_{q,d_0} \alpha^{2d_0} \eta^{2d_0-1} |B|^{2d_0-1}, \quad c_{q,d_0} = \left(8q \frac{q^{2d_0-2} - 1}{q - 1} \right)^{-(2d_0-1)}.$$

At order eight over $\mathbb{F}_2$, $s = 15$ and $S_{2,15} = 16383$. The large-set theorem has the explicit form

$$|B| \geq \frac{262128}{\alpha \eta} \implies T_8(B') \geq 2^{-46} \alpha^{16} \eta^{15} |B|^{15}.$$

The respected proportion is at least $1/(1 + \eta)$. Without a size condition the coefficient 262128^{-15} is valid. The value 2^{-270} is a slightly weaker simple power of two, since $262128 < 2^{18}$.

The conversion from graph quadruple energy to the order-eight input is an established Fourier-moment argument, following Gowers's Lemma 9.2 [1] and retaining the exact second moment as in [136]. We include the short proof to make the specialization self-contained.

Lemma II.120.7 (Graph-energy interpolation). *Let B be a nonempty finite subset of a finite abelian group and let ϕ take values in another finite abelian group. For an integer $j \geq 1$, let $G_j(B, \phi)$ count the ordered additive $2j$ -tuples in B that are respected by ϕ . If $0 < \gamma \leq 1$ and $G_2(B, \phi) \geq \gamma |B|^3$, then*

$$G_8(B, \phi) \geq \gamma^7 |B|^{15}.$$

Proof. Let Γ be the graph of ϕ , let $m = |\Gamma| = |B|$, and write H for the product of the two ambient groups and $N = |H|$. For $\chi \in \hat{H}$, define the unnormalized character sum

$$F_\Gamma(\chi) = \sum_{z \in \Gamma} \overline{\chi(z)}.$$

Orthogonality and Parseval give

$$G_j(B, \phi) = \frac{1}{N} \sum_{\chi \in \hat{H}} |F_\Gamma(\chi)|^{2j}, \quad \frac{1}{N} \sum_{\chi \in \hat{H}} |F_\Gamma(\chi)|^2 = m.$$

Consequently $p_\chi = |F_\Gamma(\chi)|^2 / (Nm)$ is a probability distribution. For the random variable $Z_\chi = |F_\Gamma(\chi)|^2$,

$$\mathbb{E}_p Z = \frac{G_2(B, \phi)}{m} \geq \gamma m^2, \quad \mathbb{E}_p Z^7 = \frac{G_8(B, \phi)}{m}.$$

Jensen's inequality for the convex function $t \mapsto t^7$ on $[0, \infty)$ yields $G_8(B, \phi)/m \geq (\gamma m^2)^7$, proving the result. $\square$

Substituting $\alpha = \gamma^7$ in the binary order-eight restriction gives

$$|B| \geq 262128 \gamma^{-7} \eta^{-1}, \quad T_8(B') \geq 2^{-46} \gamma^{112} \eta^{15} |B|^{15}.$$

These bounds combine the established interpolation with the new all-tuples restriction estimate.

For $k \geq 1$, an additive k -cube has 2^k entries and a translation-invariant relation space of dimension $2^k - k - 1$, leaving $s = k + 1$ free coordinates. Thus if at least $\alpha|B|^{k+1}$ cubes are respected by ϕ , the new theorem retains at least

$$\frac{\alpha^{k+2}\eta^{k+1}}{2(4q)^{k+1}}|B|^{k+1}$$

cubes with respected proportion at least $1/(1 + \eta)$ under $|B| \geq 8q(q^k - 1)/((q - 1)\alpha\eta)$. For cubes, “respected” means that the full system of affine-cube relations also holds for the images. Requiring only the alternating top-degree sum to vanish is a different condition. The affine-coordinate matrix of the cube vertices has rank $k + 1$ over every field: its columns at the zero vertex and the k unit vertices form an invertible affine-coordinate basis, including in characteristic two.

II.121 Sources 28 and 30 at Section 10: logarithmic selection and fixed radii [write]

Added in the sixth write, 6 October 2026.

The next fourteen sections return to Gowers’s Section 10, the local Freiman model on a Bohr neighbourhood. *Logically they continue source 12’s chapter* (Section II.33) *and source 07’s* (Section II.11), which stay the spines for paths, extension, synchronization and smoothing; they are printed here because no number printed earlier may move. Sections II.122–II.127 are source 28’s, Sections II.128–II.135 source 30’s; their introductions are Sections F.33 and F.35.

What is proved, and what is new in this report. Source 28 replaces the component selection by random windows in the logarithms of the fibre size and the anchored energy weight, and obtains a local model on $\text{Bohr}(\Gamma, \zeta_\alpha)$ from the original-domain spectrum, retaining $\alpha p(\alpha)/20$ of the domain, $p(\alpha) = (\log 2)^2/(\log(32/\alpha) \log(32/\alpha^2))$, with $|\Gamma| \leq 2^{34}(1 + \alpha)/((\log 2)\alpha^3)$, radius $\Omega(\alpha^5)$, input defect 2^{-21} and order seven (Theorem F.33.1); against the corrected formal statement ($\alpha^6/20000$, $O(\alpha^{-10})$, 2^{-43}) and source 12 ($\alpha^5/20000$, radius α^{19}). New are the logarithmic selection (Lemma II.124.1), the energy-weighted boundary estimate and the parameter chain; the coupled paths, dense extension and smoothing are credited to sources 07 and 12. Source 30 proves a fixed-radius translation estimate $2m b_B(d) \leq |B(3\rho/2)| - |B| \leq (3^r - 1)|B|$ with both constituents sharp (Theorems II.129.3, II.130.2), a margin-preserving radius $\rho/(8(3^r - 1))$ for Corollary 10.11 at the given radius (Corollary II.132.1), the 5/14 compatibility count, and a Chinese-remainder obstruction: in cyclic groups of composite order the dependence on the rank must be exponential (Corollary II.135.5). Its tree budget and compatibility inequality are source 12’s, and its extension budget a weaker form of source 12’s sharp theorem (notes there); composing its escape inequality with source 12’s theorem enlarges its extension radius (note after Corollary II.132.2). Source 12’s radius $\rho/(8d)$ is larger but selected in $(\rho/2, \rho]$; the obstruction explains the price of a fixed radius. Not to be confused: `Proofs07BohrHom.lean` uses a related fourth-moment bound for Section 7, not Section 10.

Relation to the formal statements (checked at HEAD 5d4718772). `lemma_10_10`, `corollary_10_11`, `lemma_10_12` and `theorem_10_13` (`Section10.lean`, lines 362, 374, 401, 473) are exact companions (`Proofs10Bohr.lean`, lines 277, 607; `Proofs10Induced.lean`, line 145; `Proofs10Main.lean`, line 1292). Nothing in this chapter is formalized; whether source 28’s theorem implies `theorem_10_13` with better parameters is a marked question (item II.47).

II.122 Conventions and finite domains

Source 28, delivered Section 2.

All groups are written additively except that characters take their values in the unit circle. Expectations over nonempty finite sets are uniform. For G of order N , set

$$\widehat{u}(\chi) = \sum_{s \in G} u(s) \overline{\chi(s)}, \quad (u * v)(s) = \sum_{t \in G} u(t) v(s - t).$$

Then

$$u(s) = \frac{1}{N} \sum_{\chi \in \widehat{G}} \widehat{u}(\chi) \chi(s), \quad \sum_{\chi} |\widehat{u}(\chi)|^2 = N \sum_s |u(s)|^2.$$

Part II explicitly switches to normalized Fourier coefficients, $\widehat{f}(\chi) = \mathbb{E}_x f(x) \overline{\chi(x)}$, denoted a_χ in its concentration argument. These conventions are never mixed inside one calculation.

Definition II.122.1 (Bohr sets). For $\Gamma \subseteq \widehat{G}$ and $u \geq 0$,

$$\text{Bohr}(\Gamma, u) = \{s \in G : \|\arg \chi(s)/(2\pi)\|_{\mathbb{R}/\mathbb{Z}} \leq u \text{ for every } \chi \in \Gamma\}.$$

Here $\|t\|_{\mathbb{R}/\mathbb{Z}}$ is distance to the nearest integer. Bohr sets are symmetric and contain 0.

Definition II.122.2 (Relations and exact local models). An order- k relation in a domain (X, r) is an ordered tuple $(x_1, \dots, x_{2k}) \in X^{2k}$ satisfying

$$\sum_{j=1}^k r(x_j) = \sum_{j=k+1}^{2k} r(x_j).$$

It is respected by $\phi : X \rightarrow H$ if the same equality holds with ϕ in place of r . A $(1-\eta)$ -homomorphism of order k respects at least that fraction of all order- k relations. Repetition of points is allowed.

A Freiman homomorphism on $A \subseteq G$ is a map $F : A \rightarrow H$ respecting all order-two relations. It need not be the restriction of a homomorphism defined on all of G .

A domain is $(B, \sigma M)$ -invariant if

$$||X_{s+d}| - |X_s|| \leq \sigma M \quad (s \in G, d \in B).$$

A *fibre-saturated* subset is a union of whole fibres X_s . For such a subset W , the statement that it survives the shift d means that most of its *points*, counted with their fibre multiplicities, have their destination fibre also inside W . We use the explicit expression

$$W^{(d)} = \{w \in W : r(w) + d \in r(W)\}.$$

This avoids treating a shift of a multifunction domain as a pointwise permutation of its points.

II.123 Weighted local errors and a useful anchor

Source 28, delivered Section 3.

Fix a nonempty symmetric $B \subseteq G$ containing 0. In this section

$$|X| = \alpha MN, \quad |X_s| \leq M, \quad 0 < \eta, \quad 0 < \sigma \leq \eta\alpha^2,$$

and $\phi : X \rightarrow H$ is a $(1 - \eta)$ -homomorphism of order two on a $(B, \sigma M)$ -invariant domain. Put $n_s = |X_s|$. For $x, y \in X$, define

$$q(x, y) = \sum_{d \in G} n_{r(x)+d} n_{r(y)+d}, \quad (\text{II.123.1})$$

$$b(x, y) = \sum_{d \in B} n_{r(x)+d} n_{r(y)+d}, \quad (\text{II.123.2})$$

$$e(x, y) = \#\{(u, v) \in X^2 : r(u) - r(x) = r(v) - r(y) \in B, \phi(u) - \phi(x) \neq \phi(v) - \phi(y)\}, \quad (\text{II.123.3})$$

$$\epsilon(x, y) = e(x, y)/b(x, y). \quad (\text{II.123.4})$$

The denominator is positive because $0 \in B$. Write

$$Q(x) = \sum_y q(x, y), \quad E(x) = \sum_y \epsilon(x, y) q(x, y), \quad R(x) = n_{r(x)}, \quad T = \sum_x Q(x).$$

Lemma II.123.1 (Energy bounds). *One has*

$$q(x, y) \leq \alpha M^2 N, \quad Q(x) \leq \alpha^2 M^3 N^2, \quad T \geq \alpha^4 M^4 N^3.$$

Proof. For a fixed difference there are at most $|X|M$ ordered point pairs: choose one point and then choose the other within a specified fibre. This proves the first inequality, and summing over y proves the second. If a_d is the number of ordered pairs of difference d , then

$$T = \sum_d a_d^2 \geq \frac{1}{N} \left(\sum_d a_d \right)^2 = \frac{|X|^4}{N}.$$

□

Lemma II.123.2 (Transfer to anchored weighted error).

$$\sum_x E(x) \leq 15\eta T.$$

Proof. This is the local error-transfer mechanism of Gowers's Lemma 10.2, written with explicit finite counts. Let $L = \eta\alpha^2 M$, and call a point thick if its fibre has size at least $5L$. For thick x, y , and any (u, v) with $r(u) - r(x) = r(v) - r(y) \in B$, invariance gives

$$\frac{b(u, v)}{b(x, y)} \leq \frac{(n_{r(x)} + 2L)(n_{r(y)} + 2L)}{(n_{r(x)} - L)(n_{r(y)} - L)} \leq \left(\frac{7}{4}\right)^2 < 4.$$

If the increments at (x, u) and (y, v) disagree, then for any pair (z, w) with $r(w) - r(z) = r(y) - r(x)$, at least one of the two equal-base-difference pairs

$$((x, y), (z, w)), \quad ((u, v), (z, w))$$

has unequal label differences. Count the resulting sextuples with weight $1/b(x, y)$. For a failure of the first kind, summing over (u, v) has total weight exactly one. For a failure of the second kind, symmetry of B and the displayed ratio bound give total weight at most four. The order-two defect therefore bounds the thick contribution by $5\eta T$.

There are at most $5\eta\alpha^2 MN$ nonthick points. Since $\epsilon \leq 1$, the contribution of pairs with at least one nonthick endpoint is, by lemma II.123.1, at most

$$2(5\eta\alpha^2 MN)(\alpha^2 M^3 N^2) = 10\eta\alpha^4 M^4 N^3 \leq 10\eta T.$$

Adding the contributions proves the result. $\square$

Lemma II.123.3 (An anchor with simultaneous energy and error control). *There exists $x \in X$ such that*

$$R(x) \geq \frac{\alpha^2 M}{2}, \quad S := Q(x) \geq \frac{\alpha^3 M^3 N^2}{4}, \quad E(x) \leq 60\eta S.$$

Proof. Let $A = \{x : R(x) \geq \alpha^2 M/2\}$. The discarded points are at most $\alpha^2 MN/2$, so

$$\sum_{x \notin A} Q(x) \leq \frac{\alpha^4 M^4 N^3}{2} \leq \frac{T}{2}.$$

Consequently

$$\sum_{x \in A} \left(Q(x) - \frac{E(x)}{60\eta} \right) \geq \frac{T}{2} - \frac{15\eta T}{60\eta} = \frac{T}{4}.$$

Some $x \in A$ has the summand at least $T/(4|A|)$, which is at least $T/(4|X|) \geq \alpha^3 M^3 N^2/4$. Its summand is positive, giving both required conclusions. $\square$

Fix this anchor. Abbreviate $q(y) = q(x, y)$ and $\epsilon(y) = \epsilon(x, y)$, and put $q_{\max} = \alpha M^2 N$. For points y, z with $r(z) - r(y) \in B$,

$$|q(z) - q(y)| \leq \sigma \alpha M^2 N = \sigma q_{\max}. \quad (\text{II.123.5})$$

Indeed, shift one of the two fibre-count factors in the defining convolution and use $\sum_s n_s = |X|$. The same estimate is valid for the convolution considered at arbitrary base coordinates.

II.124 Logarithmic component selection

Source 28, delivered Section 4.

The next lemma is the central refinement. Additive windows have to be small enough for the smallest fibres they can encounter. Logarithmic windows encode the desired ratio bound directly. Moreover, the boundary is measured by q , the same weight as the useful energy. The resulting comparison is local to the selected window.

Lemma II.124.1 (Selection with logarithmic windows). *Under the anchored hypotheses above, assume*

$$0 < \sigma \leq \frac{\eta(\log 2)\alpha^2}{2048(1 + \alpha)}, \quad 0 < \eta \leq 2^{-21}. \quad (\text{II.124.1})$$

There is a nonempty fibre-saturated $W \subseteq X$ with

$$|W| \geq \frac{p(\alpha)\alpha}{16}|X|, \quad (\text{II.124.2})$$

$$\frac{\max_W q}{\min_W q} \leq 2, \quad \frac{\max_W R}{\min_W R} \leq 2, \quad (\text{II.124.3})$$

$$\min_W q \geq \frac{\alpha q_{\max}}{32}, \quad \min_W R \geq \frac{\alpha^2 M}{32}, \quad (\text{II.124.4})$$

$$\mathbb{E}_{w \in W} \epsilon(w) \leq 960\eta, \quad (\text{II.124.5})$$

$$|W^{(d)}| \geq (1 - \eta)|W| \quad (d \in B). \quad (\text{II.124.6})$$

Proof. Step 1: discard two low-energy tails. Put

$$a_0 = \alpha/16, \quad b_0 = \alpha^2/16, \quad X_* = \{y : q(y)/q_{\max} \geq a_0, R(y)/M \geq b_0\}.$$

Discarding points with $q(y) < a_0 q_{\max}$ loses at most

$$a_0 q_{\max} |X| = \frac{\alpha^3 M^3 N^2}{16}$$

of the q -weight. Discarding the other small fibres loses at most

$$(b_0 M N) q_{\max} = \frac{\alpha^3 M^3 N^2}{16}.$$

Both estimates are at most $S/4$. Thus

$$S_* := \sum_{y \in X_*} q(y) \geq S/2. \quad (\text{II.124.7})$$

Step 2: random multiplicative windows. Set $h = (\log 2)/2$. Choose u, v independently, uniformly on

$$I_q = [\log a_0 - h, h], \quad I_R = [\log b_0 - h, h].$$

The random set $W_{u,v}$ consists of the points y satisfying

$$|\log(q(y)/q_{\max}) - u| \leq h, \quad |\log(R(y)/M) - v| \leq h.$$

All quantities inside the logarithms are positive. For a fixed point of X_* , the admissible center interval in each coordinate has length $2h$. Hence its inclusion probability is exactly

$$p = \frac{(2h)^2}{|I_q||I_R|} = \frac{(\log 2)^2}{\log(32/\alpha) \log(32/\alpha^2)}.$$

Every other point has inclusion probability at most p . Any nonempty selected window automatically satisfies equations (II.124.3) and (II.124.4): the ratio inside either coordinate interval is $\exp(2h) = 2$, while its lowest possible endpoint is respectively $a_0/2$ or $b_0/2$.

Step 3: a boundary inside the selected window. Define

$$\Delta_q = \frac{64\sigma}{\alpha}, \quad \Delta_R = \frac{64\sigma}{\alpha^2}, \quad \beta = \frac{\Delta_q + \Delta_R}{h} = \frac{128\sigma(1 + \alpha)}{\alpha^2 \log 2}.$$

The hypotheses imply $\sigma \leq \alpha^2/128$ and $0 < \Delta_q, \Delta_R < h$. Let $V_{u,v}$ be $W_{u,v}$ minus the inner window obtained by replacing the half-widths h with $h - \Delta_q, h - \Delta_R$. For every fixed point of X , even one outside X_* , the probability of lying in $V_{u,v}$ is at most $p\beta$. To see this, a boundary interval in the first center coordinate has total length at most $2\Delta_q$, and the other coordinate has length at most $2h$; interchange the coordinates and use the union bound. Truncation by I_q, I_R only decreases these lengths.

This choice of boundary has a deterministic consequence. If $y \in W_{u,v}$ and $d \in B$, invariance first gives

$$n_{r(y)+d} \geq R(y) - \sigma M \geq \alpha^2 M/32 - \sigma M > 0.$$

Choose z in that fibre. By equation (II.123.5) and the lower bounds in equation (II.124.4),

$$\begin{aligned} \left| \log \frac{q(z)}{q(y)} \right| &\leq \frac{\sigma}{\alpha/32 - \sigma} \leq \frac{64\sigma}{\alpha} = \Delta_q, \\ \left| \log \frac{R(z)}{R(y)} \right| &\leq \frac{\sigma}{\alpha^2/32 - \sigma} \leq \frac{64\sigma}{\alpha^2} = \Delta_R. \end{aligned}$$

We used $|\log a - \log b| \leq |a - b|/\min(a, b)$. Thus every point in $W_{u,v} \setminus V_{u,v}$ has its shifted fibre entirely in $W_{u,v}$, simultaneously for every $d \in B$.

Step 4: one correlated selection inequality. For each pair of centers write

$$S_{u,v} = \sum_{X_* \cap W_{u,v}} q, \quad E_{u,v} = \sum_{W_{u,v}} \epsilon q, \quad Q_{V,u,v} = \sum_{V_{u,v}} q.$$

The probability estimates and the anchor inequalities give

$$\mathbb{E} S_{u,v} \geq pS/2, \quad \mathbb{E} E_{u,v} \leq 60\eta pS, \quad \mathbb{E} Q_{V,u,v} \leq p\beta S.$$

Therefore

$$\mathbb{E} \left[S_{u,v} - \frac{E_{u,v}}{480\eta} - \frac{Q_{V,u,v}}{8\beta} \right] \geq pS/4. \quad (\text{II.124.8})$$

Choose centers for which the bracket is at least $pS/4$. Set $W = W_{u,v}$, $V = V_{u,v}$, and $S_W = S_{u,v}$. Then

$$S_W \geq pS/4, \quad E_{u,v} \leq 480\eta S_W, \quad Q_{V,u,v} \leq 8\beta S_W.$$

In particular $W \neq \emptyset$, and

$$|W| \geq \frac{S_W}{q_{\max}} \geq \frac{p\alpha^2 MN}{16},$$

which is equation (II.124.2).

Step 5: convert weighted errors only after selecting. Since the ratio of the largest and smallest q -weights in W is at most two,

$$\mathbb{E}_W \epsilon \leq \frac{\sum_W \epsilon q}{|W| \min_W q} \leq 480\eta \frac{\max_W q}{\min_W q} \leq 960\eta.$$

Similarly, using $S_W \leq \sum_W q \leq |W| \max_W q$,

$$\frac{|V|}{|W|} \leq \frac{Q_{V,u,v}}{|W| \min_W q} \leq 16\beta \leq \eta.$$

Step 3 now implies equation (II.124.6). □

II.124.1 Why the density powers disappear

The original additive-window estimate separately replaces the selected energy by its global upper bound and the selected cardinality by its global lower bound. That loses their correlation. Localizing the boundary to the selected window removes another unnecessary loss. Weighting the boundary by q then makes it comparable directly with the selected energy. Finally, multiplicative windows need only sample a logarithmic range in each coordinate.

The two logarithmic factors have a precise origin: one pays for choosing a multiplicative scale of the anchored energy, and the other for choosing a multiplicative scale of the fibre size. All subsequent error bounds are proportional to the mass in the selected window. They introduce no additional inverse power of α . No optimality of the density powers, logarithms, or numerical constants is asserted.

II.125 From an averaged anchor error to exact local differences

Source 28, delivered Section 5.

The previous transport report already developed coupled paths and collision-to-mode exactification [133]. Its collision input is a normalized directional law; the comparison with a raw additive-energy statistic is explicitly a separate requirement. We supply that comparison for the anchored error furnished by the component selection, then carry the average directly through the final restriction. This avoids imposing the nested exceptional-set conclusions of the original Lemma 10.6.

Throughout this section, G is a finite abelian group, H is an abelian group, U is a finite set with a map $r : U \rightarrow G$, and $F : U \rightarrow H$. Write $U_s = r^{-1}(s)$ and $R(u) = |U_{r(u)}|$. The target group need not be finite. Every probability distribution used below nevertheless has finite support.

Lemma II.125.1 (Anchored agreement selects a mode). *Let $B \subseteq G$ be nonempty, $W \subseteq U$ nonempty, and $x \in U$. Assume that $U_{r(x)+d}$ and $U_{r(w)+d}$ are nonempty for every $d \in B$ and $w \in W$. Put*

$$b_d(w) = |U_{r(x)+d}| |U_{r(w)+d}|,$$

and let $p(d, w)$ be the probability that

$$F(y) - F(x) \neq F(z) - F(w),$$

where y, z are independent uniform elements of $U_{r(x)+d}, U_{r(w)+d}$, respectively. Define

$$\epsilon(w) = \frac{\sum_{d \in B} b_d(w) p(d, w)}{\sum_{d \in B} b_d(w)}.$$

Suppose that

$$\max_{d \in B} b_d(w) \leq 4 \min_{d \in B} b_d(w) \quad (w \in W), \quad \mathbb{E}_{w \in W} \epsilon(w) \leq 1024\eta.$$

If $0 < \eta \leq 2^{-21}$, there are $D \subseteq B$ and $\psi : D \rightarrow H$ such that

$$|D| \geq \frac{15}{16} |B|$$

and, for every $d \in D$,

$$\mathbb{E}_{w \in W} e_d(w) \leq \frac{1}{31}, \quad e_d(w) = \mathbb{P}_{z \in U_{r(w)+d}} (F(z) - F(w) \neq \psi(d)). \quad (\text{II.125.1})$$

Proof. The weight comparison implies $\mathbb{E}_{d \in B} p(d, w) \leq 4\epsilon(w)$, hence

$$\mathbb{E}_{d \in B, w \in W} p(d, w) \leq 4096\eta.$$

Let

$$D = \left\{ d \in B : \mathbb{E}_{w \in W} p(d, w) \leq \frac{1}{32} \right\}.$$

Markov's inequality gives

$$\frac{|B \setminus D|}{|B|} \leq 2^{17}\eta \leq \frac{1}{16}.$$

For $d \in D$, let A_d be the distribution of $F(y) - F(x)$ for uniform $y \in U_{r(x)+d}$. Let Q_d be the distribution of $F(z) - F(w)$ obtained by choosing w uniformly in W and then z uniformly in $U_{r(w)+d}$. Independent samples from A_d, Q_d agree with probability at least $31/32$. Their agreement probability is no larger than the greatest atom of A_d . That distribution therefore has a unique atom $\psi(d)$ of mass at least $31/32$.

The event that the first sample equals $\psi(d)$ and the second does not is contained in the disagreement event. Independence gives

$$\frac{31}{32} Q_d(H \setminus \{\psi(d)\}) \leq \frac{1}{32}.$$

The probability on the left is precisely $\mathbb{E}_{w \in W} e_d(w)$, proving the assertion. $\square$

Lemma II.125.2 (Two paths with a common endpoint). *Assume W is a nonempty union of whole fibers and its positive fiber sizes vary by a factor of at most two. Let $D \subseteq G$ and suppose that $U_{r(w)+d}$ is nonempty whenever $w \in W, d \in D$. Suppose also that*

$$|\{w \in W : r(w) + d \in r(W)\}| \geq (1 - \eta)|W| \quad (d \in D). \quad (\text{II.125.2})$$

If $\psi : D \rightarrow H$ satisfies $\mathbb{E}_{w \in W} e_d(w) \leq e$ for every $d \in D$, with e_d defined as in (II.125.1), and

$$6e + 2\eta < 1,$$

then ψ is a Freiman homomorphism of order two.

Proof. Choose w uniformly in W and z uniformly in $U_{r(w)+a}$, where $a \in D$. The probability that $z \notin W$ is at most η . For any particular $z \in W$, its probability under this sampling law is

$$\frac{\mathbf{1}_{\{r(z)-a \in r(W)\}}}{|W|} \frac{|U_{r(z)-a}|}{|U_{r(z)}|} \leq \frac{2}{|W|}.$$

Consequently every nonnegative function h on W obeys

$$\mathbb{E}[\mathbf{1}_{\{z \in W\}} h(z)] \leq 2\mathbb{E}_{z \in W} h(z). \quad (\text{II.125.3})$$

Fix $a, b, c, d \in D$ with $a + b = c + d = s$. Choose one common starting point w uniformly in W and one common endpoint u uniformly in $U_{r(w)+s}$. If this endpoint fiber is empty, use a dummy point and declare both paths unsuccessful. Choose intermediates independently from $U_{r(w)+a}$ and $U_{r(w)+c}$.

For the path with increments (a, b) , the first edge fails with probability at most e , and the intermediate lies outside W with probability at most η . When the intermediate is in W , the endpoint fiber is nonempty by hypothesis. The second edge fails with probability at most $2e$, by

(II.125.3). Thus the total failure probability is at most $3e + \eta$. The other path has the same bound. Since $6e + 2\eta < 1$, both paths succeed with positive probability. On that event,

$$\psi(a) + \psi(b) = F(u) - F(w) = \psi(c) + \psi(d).$$

This proves the required exact relation. The common endpoint, rather than two separately sampled endpoint values, is what makes the conclusion valid. $\square$

Corollary II.125.3 (Exact modes on a dense set of directions). *Under the hypotheses of Lemma II.125.1, assume in addition that W is a union of whole fibers with size ratio at most two and that (II.125.2) holds for every $d \in B$. Then the map $\psi : D \rightarrow H$ produced there is a Freiman homomorphism of order two.*

Proof. Apply Lemma II.125.2 with $e = 1/31$, since $6/31 + 2^{-20} < 1$. $\square$

II.125.1 Regular Bohr sets and dense difference extension

We recall the one-scale argument and extension criterion from the prior Bohr report [134], in the exact form used here. For a finite set K of characters write

$$B(K, \rho) = \{g \in G : \|\arg \chi(g)/(2\pi)\|_{\mathbb{R}/\mathbb{Z}} \leq \rho \text{ for every } \chi \in K\}.$$

Lemma II.125.4 (One selected radius). *Let $1 \leq k = |K| \leq R$ and $0 < \rho_0 \leq 1/2$. There exists $\rho \in [\rho_0/2, \rho_0]$ such that, with*

$$B = B(K, \rho), \quad C = B\left(K, \frac{\rho_0}{6R}\right), \quad \beta = 4^{-1/3} > \frac{5}{8},$$

one has $|B \cap (B + t)| \geq \beta|B|$ for every $t \in C$.

Proof. Partition each coordinate interval $[-\rho_0, \rho_0]$ into four intervals of length $\rho_0/2$. Any two group points in the same product cell differ by an element of $B(K, \rho_0/2)$. Translating each nonempty cell by one of its points gives

$$|B(K, \rho_0)| \leq 4^k |B(K, \rho_0/2)|.$$

Set $\tau = \rho_0/(6k)$. Among the $3k$ adjacent steps from $\rho_0/2$ to ρ_0 , some step has cardinality ratio at most $4^{1/3}$. At its upper endpoint ρ ,

$$|B(K, \rho - \tau)| \geq \beta |B(K, \rho)|.$$

For $t \in C \subseteq B(K, \tau)$, the inner set is contained in $B \cap (B + t)$, proving the result. $\square$

The radius in Lemma II.125.4 is selected before constructing the anchor and W . The anchor errors are then defined using this actual B ; no subsequent restriction of their averaging domain is needed.

Lemma II.125.5 (Dense Freiman extension). *Let $D \subseteq B \subseteq G$, with $B \neq \emptyset$, satisfy $|B \setminus D| \leq |B|/16$. Suppose $|B \cap (B + t)| \geq \beta|B|$ for $t \in C$, where $\beta > 5/8$, and $\psi : D \rightarrow H$ is a Freiman homomorphism of order two. Then $C \subseteq D - D$, and the difference map induced by ψ restricts to a Freiman homomorphism $\Psi : C \rightarrow H$ of order two. If $0 \in C$, then $\Psi(0) = 0$.*

Proof. The Freiman property makes $\Delta(a - b) = \psi(a) - \psi(b)$ well defined on $D - D$. For $t \in C$,

$$|D \cap (D - t)| \geq \left(\beta - \frac{1}{8}\right) |B| > \frac{|B|}{2}. \quad (\text{II.125.4})$$

In particular $C \subseteq D - D$, so define $\Psi = \Delta|_C$.

For $c, d \in C$, the intersection $D \cap (D - c) \cap (D - c - d)$ has size at least

$$\left(1 - \frac{3}{16} - 2(1 - \beta)\right) |B| > \frac{|B|}{16}.$$

Indeed the chain of ambient translates $B, B - c, B - c - d$ loses at most the two boundary sizes for c and d , and the three translated hole sets lose at most $3|B|/16$. A point in the intersection proves that $c + d \in D - D$ and

$$\Delta(c + d) = \Psi(c) + \Psi(d).$$

This statement holds even when $c + d \notin C$. Two equal pair-sums from C therefore have equal Ψ pair-sums. Also $\Delta(0) = 0$. $\square$

Lemma II.125.6 (Averaged rows give exact compatibility). *Under the hypotheses of Lemma II.125.5, suppose W is nonempty, all fibers $U_{r(w)+d}$ for $w \in W, d \in D$ are nonempty, and (II.125.1) holds. Then there is $W_1 \subseteq W$ with*

$$|W_1| \geq \frac{27}{31} |W| > \frac{5}{6} |W|$$

such that

$$F(w_1) - F(w_2) = \Psi(r(w_1) - r(w_2))$$

whenever $w_1, w_2 \in W_1$ and $r(w_1) - r(w_2) \in C$.

Proof. Set

$$W_1 = \left\{ w \in W : \mathbb{E}_{d \in D} e_d(w) \leq \frac{1}{4} \right\}.$$

The average row error is at most $1/31$, so Markov's inequality gives the stated size. Fix $w_1, w_2 \in W_1$ and put $c = r(w_1) - r(w_2) \in C$. For $d \in I = D \cap (D - c)$, the target fibers $U_{r(w_1)+d}$ and $U_{r(w_2)+d+c}$ coincide. Moreover,

$$\sum_{d \in I} (e_d(w_1) + e_{d+c}(w_2)) \leq \frac{|D|}{2} < |I|,$$

where the strict inequality follows from (II.125.4). For some $d \in I$ the two conditional errors therefore sum to less than one. Their good subsets of the common nonempty fiber intersect. At a point z in the intersection,

$$F(z) - F(w_1) = \psi(d), \quad F(z) - F(w_2) = \psi(d + c).$$

Subtracting proves the assertion. $\square$

Remark II.125.7 (The auxiliary invariance requirement). Suppose the ambient fibers change by at most σM_U under directions in B , the anchor fiber has size at least $\alpha^2 M_U/2$, and the fibers on W have size at least $\alpha^2 M_U/32$. The nonemptiness and fourfold weight comparison in Lemma II.125.1 follow from $\sigma \leq \alpha^2/96$: each translated fiber lies between two thirds and four thirds of its original size. The smaller auxiliary condition appearing in the original Lemma 10.6 is therefore unnecessary for this replacement. The component theorem supplies the averaged anchor error, and the present lemmas finish the exactification directly.

II.125.2 A higher-order conclusion at the same parameters

The full ambient invariance in the present application permits the prefix-sensitive criterion of the prior transport report [133]. We give the short specialization, because it strengthens the output without changing the spectrum, radius, or retained set.

Lemma II.125.8 (Ambient invariance and longer paths). *In the setting of Lemma II.125.2, write $R_s = |U_s|$ and assume $R(w) \geq m > 0$ for $w \in W$. Suppose in addition that*

$$|R_{s+d} - R_s| \leq L \quad (s \in G, d \in D),$$

and put $\tau = L/m$. For any integer $k \geq 2$,

$$(2k + \tau k(k-1))(e + \eta) < 1$$

implies that ψ is a Freiman homomorphism of order k .

Proof. Use the unnormalized subprobability law of the paths that have survived each prefix; never condition the law to have total mass one. At a prefix of j directions, the mass of a point $z \in W$ is at most

$$\frac{1}{|W|} \frac{R_{r(z)-s_j}}{R_{r(z)}} \leq \frac{1+j\tau}{|W|},$$

where s_j is the sum of those directions. The last inequality follows by applying the ambient invariance bound j times. Stop a path when it leaves W or an edge has the wrong label. For a uniformly sampled source in W , the probability of either event on a fixed direction is at most $e + \eta$. The marginal domination at prefix j therefore bounds the failure probability of a length- k path by

$$\sum_{j=0}^{k-1} (1+j\tau)(e + \eta) = \left(k + \frac{\tau k(k-1)}{2} \right) (e + \eta).$$

Compare two length- k paths with the same initial and terminal indices, using a common sampled initial point and a common sampled endpoint as in Lemma II.125.2. The displayed hypothesis ensures simultaneous success with positive probability. Their telescoping label sums agree, which proves the order- k Freiman property. $\square$

Lemma II.125.9 (Order of the induced difference map). *If D is nonempty and $\psi : D \rightarrow H$ is a Freiman homomorphism of order $m \geq 2$, its induced difference map on $D - D$ is a Freiman homomorphism of order $\lfloor m/2 \rfloor$ and takes the value zero at zero.*

Proof. Order m implies every lower order by padding the two sides with copies of one fixed point of D . Order two makes the difference map well defined. A relation between two sums of q differences from $D - D$ expands, after moving their negative terms to the opposite side, into a relation between two sums of $2q$ points of D . Applying the order- $2q$ property proves the conclusion whenever $2q \leq m$. $\square$

Corollary II.125.10 (Order seven on the Bohr neighborhood). *Under the parameters of the main construction,*

$$\eta = 2^{-21}, \quad e = \frac{1}{31}, \quad m = \frac{\alpha^2 M_U}{32}, \quad L = \sigma M_U, \quad \sigma = \frac{(\log 2)\alpha^2}{2^{32}(1+\alpha)},$$

the map ψ is a Freiman homomorphism of order fifteen on D . Its difference map is a Freiman homomorphism of order seven on $D - D$; in particular, so is the final map Ψ on C .

Proof. Here $\tau = 32\sigma/\alpha^2 < 2^{-27}$ and

$$\left(30 + 210 \cdot 2^{-27}\right) \left(\frac{1}{31} + 2^{-21}\right) < 1.$$

The last inequality is a rational comparison. Apply Lemma II.125.8 with $k = 15$, followed by Lemma II.125.9. $\square$

II.126 Fourier smoothing and the proof of the main theorem

Source 28, delivered Section 6.

We now apply the component and repair lemmas to a fourfold difference domain. This is where the order-eight assumption is used. The Fourier smoothing estimate itself is the full-moment refinement already present in source 12 [134].

Lemma II.126.1 (Fourfold smoothing with the original spectrum). *Let $M > 0$, $0 < \alpha \leq 1$, and $0 \leq g(s) \leq M$ on G , with $\sum_s g(s) = \alpha MN$, and put*

$$h = g * \tilde{g} * g * \tilde{g}, \quad \tilde{g}(s) = g(-s), \quad L = \alpha^3 M^4 N^3.$$

For $\lambda > 0$, let $\Gamma = \{\chi : |\hat{g}(\chi)| \geq \lambda MN\}$. Then

$$|\Gamma| \leq \alpha/\lambda^2, \quad \sum_{\chi} |\hat{g}(\chi)|^4 \leq \alpha^3 M^4 N^4.$$

If $d \in \text{Bohr}(\Gamma, \rho_0)$, then

$$|h(s+d) - h(s)| \leq \left(2\pi\rho_0 + \frac{2\lambda^2}{\alpha^2}\right) L \quad (s \in G). \quad (\text{II.126.1})$$

Proof. Since $g \geq 0$, $|\hat{g}(\chi)| \leq \alpha MN$. Also Parseval and $g^2 \leq Mg$ give

$$\sum_{\chi} |\hat{g}(\chi)|^2 \leq \alpha M^2 N^2.$$

The cardinality bound and the fourth-moment bound follow immediately. For the small-spectrum tail,

$$\sum_{\chi \notin \Gamma} |\hat{g}(\chi)|^4 \leq \lambda^2 M^2 N^2 \sum_{\chi} |\hat{g}(\chi)|^2 \leq \alpha \lambda^2 M^4 N^4.$$

Fourier inversion gives

$$h(s+d) - h(s) = \frac{1}{N} \sum_{\chi} |\hat{g}(\chi)|^4 \chi(s)(\chi(d) - 1).$$

Use $|\chi(d) - 1| \leq 2\pi\rho_0$ on Γ , and at most two on its complement. Division by L yields equation (II.126.1). $\square$

Proof of theorem F.33.1. Set

$$\eta_0 = 2^{-21}, \quad \sigma = \frac{(\log 2)\alpha^2}{2^{32}(1+\alpha)}, \quad \lambda^2 = \frac{\sigma\alpha^2}{4}, \quad \rho_0 = \frac{\sigma}{4\pi}. \quad (\text{II.126.2})$$

Thus $\sigma = \eta_0(\log 2)\alpha^2/[2048(1 + \alpha)]$. In particular

$$\sigma \leq \eta_0\alpha^2, \quad \sigma < \alpha^2/128, \quad 0 < \rho_0 < 1/2.$$

The threshold λ in equation (II.126.2) is exactly equation (F.33.2). Moreover

$$R_\alpha = \frac{\alpha}{\lambda^2} = \frac{4}{\sigma\alpha}, \quad \frac{\rho_0}{6R_\alpha} = \frac{\sigma^2\alpha}{96\pi} = \zeta_\alpha.$$

The trivial character is in Γ , since $\lambda < \alpha$. Thus $1 \leq |\Gamma| \leq R_\alpha$.

Form the finite domain

$$U = X^4, \quad r_U(x_1, x_2, x_3, x_4) = r(x_1) + r(x_2) - r(x_3) - r(x_4),$$

and define $F : U \rightarrow H$ by the same signed expression in ϕ . Its cardinality and fibre bound are

$$|U| = |X|^4 = \alpha LN, \quad |U_s| \leq L, \quad L = \alpha^3 M^4 N^3.$$

For the bound on a fibre, choose any three of the four points and then choose the fourth in one specified fibre of X . Its fibre-count function is the h of lemma II.126.1.

An order-two relation on U expands into an order-eight relation on X . More precisely, reorder its sixteen independent coordinates to put its eight positive and eight negative occurrences on opposite sides. This is a bijection between the two sets of tested tuples and preserves success and failure. Consequently F is a $(1 - \eta_0)$ -homomorphism of order two.

By lemma II.126.1 and equation (II.126.2), the domain U is $(\text{Bohr}(\Gamma, \rho_0), \sigma L)$ -invariant. Apply lemma II.125.4 now, before selecting any anchor, to obtain

$$B = \text{Bohr}(\Gamma, \rho), \quad \rho \in [\rho_0/2, \rho_0], \quad |B \cap (B + c)| \geq 4^{-1/3}|B| \quad (c \in C).$$

The same invariance holds on the smaller B . Apply lemmas II.123.2, II.123.3 and II.124.1 to U , using L in place of M . They give an anchor $x \in U$ and a component $W \subseteq U$ with

$$|W| \geq \frac{p(\alpha)\alpha}{16}|U|, \quad \mathbb{E}_W \epsilon \leq 960\eta_0, \quad R(w) \geq \alpha^2 L/32, \quad R(x) \geq \alpha^2 L/2.$$

The positive fibre sizes on W vary by at most a factor of two, and every shift in B preserves at least $1 - \eta_0$ of its points.

The shifted anchor and component fibres are nonempty. Indeed $\sigma \leq \alpha^2/128$ makes each such fibre size lie between $3/4$ and $5/4$ of its original size. For every $w \in W$, the largest and smallest products

$$|U_{r_U(x)+d}| |U_{r_U(w)+d}| \quad (d \in B)$$

have ratio at most $(5/3)^2 < 4$. All hypotheses of lemma II.125.1 and corollary II.125.3 follow. We obtain $D \subseteq B$, $|D| \geq 15|B|/16$, and an exact Freiman map $\psi : D \rightarrow H$ with average directional edge error at most $1/31$.

The extension and row-selection lemmas lemmas II.125.5 and II.125.6 give a map $\Psi : C \rightarrow H$ and $W_1 \subseteq W$ satisfying

$$|W_1| \geq \frac{27}{31}|W| > \frac{5p(\alpha)\alpha}{96}|U| > \frac{p(\alpha)\alpha}{20}|U|.$$

The label difference of any two points of W_1 whose base difference lies in C is given by Ψ . The stronger order-seven conclusion follows from corollary II.125.10; its hypotheses are precisely the ambient invariance and lower fibre size just established.

Finally average the cardinality of

$$Y_{x_2, x_3, x_4} = \{y \in X : (y, x_2, x_3, x_4) \in W_1\}$$

over X^3 . Some slice has at least $|W_1|/|X|^3 \geq p(\alpha)\alpha|X|/20$ points. Fix it and call it Y . The three fixed coordinates cancel from both the base difference and the label difference. This gives equation (F.33.7), and $\Psi(0) = 0$ follows from its construction as a difference map. $\square$

II.126.1 Keeping the older spectrum exactly

Some later arguments may already be organized around the old threshold. There is no need to replace that set of frequencies in order to use the new local model.

Corollary II.126.2 (Compatibility with the original threshold). *In the setting of theorem F.33.1, let*

$$\Gamma_{\text{old}} = \{\chi : |\widehat{g}(\chi)| \geq 2^{-37}\alpha^{11/2}MN\}.$$

Then the same retained set Y has an order-seven local difference model on $\text{Bohr}(\Gamma_{\text{old}}, \zeta_\alpha)$.

Proof. The ratio of the squares of the new and old thresholds is

$$\frac{2^{40} \log 2}{(1 + \alpha)\alpha^7} > 1 \quad (0 < \alpha \leq 1).$$

Thus $\Gamma \subseteq \Gamma_{\text{old}}$ and the indicated Bohr set is contained in C . Restrict Ψ . $\square$

[write] **6 October 2026; sources 28–33.** *Against the formal statement.* The intake observed that this corollary, with Theorem F.33.1, appears to give the conclusion of the formal `theorem_10_13` with the formal spectrum `domainLargeSpectrum` at the threshold λMN , a larger radius and a larger retained set, provided the order-eight defect 2^{-21} here is measured as in the formal `DomainApproxHomOfOrder`. That normalization was not checked at the intake or at the write, so the implication is recorded as a marked question (item II.47), not as a result; `theorem_10_13` is already an exact companion, so nothing formal depends on it.

The small frequency-count bound belongs to the new spectrum Γ . When the larger Γ_{old} is retained for an existing interface, the new radius and mass conclusions remain true, but its frequency count should not be relabelled $O(\alpha^{-3})$.

II.127 Consequences, comparisons, and structural limits

Source 28, delivered Section 7.

II.127.1 The retained fraction and the volume estimate

As $\alpha \downarrow 0$,

$$\frac{\alpha p(\alpha)}{20} \sim \frac{(\log 2)^2}{40} \frac{\alpha}{(\log(1/\alpha))^2}. \quad (\text{II.127.1})$$

This is an asymptotic formula for our *guaranteed lower bound*, not an asymptotic formula for the largest possible structured subset. The elementary Bohr volume estimate also becomes substantially stronger.

Lemma II.127.1 (Bohr volume). *For $0 < u \leq 1$ and $K \subseteq \widehat{G}$,*

$$|\text{Bohr}(K, u)| \geq (u/2)^{|K|} N.$$

Proof. Partition each circle coordinate into $\lceil 1/u \rceil$ intervals of length at most u . Some product cell contains at least $N/\lceil 1/u \rceil^{|K|}$ group points. Translate those points by one of them. Their differences lie in $\text{Bohr}(K, u)$, and $\lceil 1/u \rceil \leq 2/u$. $\square$

Consequently

$$|C| \geq N \exp \left[-R_\alpha \log \left(\frac{2}{\zeta_\alpha} \right) \right] = N \exp \left[-O \left(\alpha^{-3} (1 + \log(1/\alpha)) \right) \right]. \quad (\text{II.127.2})$$

This lower bound may of course be below one at a given finite N ; the actual set always contains zero. The formula records the density dependence needed for later asymptotic estimates.




Figure II.127.1: Explicit sufficient bounds. The left panel compares retained fractions with the prior fixed-spectrum source 12 bound. The right compares the corresponding polynomial radii. Both axes are logarithmic. These curves compare proved guarantees, not observed optimal outputs.

II.127.2 Local arithmetic progressions

Corollary II.127.2 (Linear difference law on short progression segments). *Let m be a positive integer and $d \in \text{Bohr}(\Gamma, \zeta_\alpha/m)$. There is $a \in H$ such that*

$$\phi(y) - \phi(z) = ja$$

whenever $y, z \in Y$ and $r(y) - r(z) = jd$, with $j \in \mathbb{Z}$ and $-m \leq j \leq m$. One can take $a = \Psi(d)$. If $G = H = \mathbb{Z}/p\mathbb{Z}$ for a prime p , this may be written as $\phi(y) - \phi(z) = c(r(y) - r(z))$ for some $c \in \mathbb{Z}/p\mathbb{Z}$.

Proof. Every jd , $|j| \leq m$, belongs to C . The order-two Freiman relation $(j-1)d + d = jd + 0$ and $\Psi(0) = 0$ give $\Psi(jd) = j\Psi(d)$, by induction; use $jd + (-jd) = 0 + 0$ for negative indices. The model

identity finishes the first assertion. In the prime cyclic case choose $c = \Psi(d)/d$ if $d \neq 0$. If $d = 0$, all the tested differences and their labels vanish. $\square$

In an arbitrary abelian group the expression ja is the correct conclusion. Division by the direction d is neither needed nor generally meaningful.

II.127.3 Several domains and a common set of directions

Corollary II.127.3 (Simultaneous local models). *Suppose t domains $X^{(i)}$ over the same G satisfy the hypotheses of theorem F.33.1, with densities $\alpha_i \geq \alpha_0 > 0$, possibly different fibre bounds and targets. Let Γ_i be their respective original-domain spectra. For*

$$C_* = \text{Bohr} \left(\bigcup_{i=1}^t \Gamma_i, \zeta_{\alpha_0} \right)$$

each domain has an order-seven exact local model on C_ , retaining at least $\alpha_0 p(\alpha_0)/20$ of its points. The union of defining frequencies has size at most tR_{α_0} .*

Proof. The functions $\alpha p(\alpha)$ and $\alpha^5/(1+\alpha)^2$ are increasing on $(0, 1]$; $(1+\alpha)/\alpha^3$ is decreasing. Apply the theorem separately and restrict each difference map to C_* . $\square$

This supplies compatible directions, not identical label maps. The distinction prevents a false inference of a single common slope across unrelated domains.

II.127.4 What is still required for a global bound

The theorem gives explicit replacements for a local Section 10 interface. A new bound for progression-free sets would additionally require an end-to-end argument propagating these parameters through the bilinear and multilinear extraction, the phase-removal step, and the density-increment iteration. The inspected formalization explicitly records outstanding obligations in those stages [131].

The new spectrum threshold is larger than the old one. This helps frequency counts, but any later theorem whose parameters were chosen using the old threshold must be reconsidered or used through corollary II.126.2. Likewise, an improved local retained fraction does not automatically repair common partitions, integer representatives of boxes, or an incorrect closing comparison. The research agenda below lists those propagation tasks as open work.

II.128 Conventions and finite translation calculus

Source 30, delivered Section 2.

All groups in the mathematical proofs are written additively. The ambient group G is finite and abelian. The target group H for partial maps may be any abelian group. For $t \in \mathbb{T}$, let

$$\|t\|_{\mathbb{T}} = \min_{n \in \mathbb{Z}} |\tilde{t} - n|,$$

where $\tilde{t}$ is any real lift. An additive character means a homomorphism $\chi : G \rightarrow \mathbb{T}$; exponentiation by $2\pi i$ converts this to the usual unit-circle convention.

Definition II.128.1 (Rectangular Bohr set). For characters $\chi = (\chi_1, \dots, \chi_r)$ and positive widths $\rho = (\rho_1, \dots, \rho_r)$, define

$$B(\chi, \rho) = \{x \in G : \|\chi_i(x)\|_{\mathbb{T}} \leq \rho_i \ (1 \leq i \leq r)\}.$$

When the characters have been fixed, we abbreviate this to $B(\rho)$, and write $B(\rho)$ for equal widths. The number r counts the listed characters, not the algebraic rank of the dual group.

Every such set contains zero and is symmetric. An empty list of characters defines the whole group; a maximum over that empty list is defined to be zero. Repeated or zero characters are allowed in the general theorem, although removing redundant constraints can improve the numerical bound. Equalities on boundaries are included throughout.

Definition II.128.2 (Translation defect and exit sets). For a nonempty subset $B \subseteq G$ and $d \in G$, set

$$b_B(d) = \text{def}_B(d) = |B| - |B \cap (B + d)|, \\ E_+(B, d) = \{x \in B : x + d \notin B\}, \quad E_-(B, d) = \{x \in B : x - d \notin B\}.$$

Lemma II.128.3 (Elementary identities). *For every finite B and d ,*

$$|E_+(B, d)| = |E_-(B, d)| = b_B(d) = b_B(-d), \quad |B \triangle (B + d)| = 2b_B(d).$$

Moreover, $b_B(d + e) \leq b_B(d) + b_B(e)$.

Proof. Translation by d bijects $B \cap (B - d)$ with $(B + d) \cap B$, so these intersections have equal size. Subtracting each size from $|B|$ proves the exit identities. Since B and $B + d$ have equal size, the two one-sided parts of their symmetric difference have equal size, giving its stated cardinality. Finally,

$$B \triangle (B + d + e) \subseteq (B \triangle (B + d)) \cup ((B + d) \triangle (B + d + e)).$$

Take cardinalities, use translation invariance, and divide by two. □

The estimates below are relative to $|B|$, not to $|G|$. This normalization is essential: an estimate that first divides an ambient boundary count by a very small Bohr density recreates the loss we seek to avoid.

II.129 The escape-tube inequality

Source 30, delivered Section 3.

II.129.1 The combinatorial mechanism

Lemma II.129.1 (Two-sided escape packing). *Let $B \subseteq G$, $d \in G$, and $m \geq 1$ be an integer. Suppose*

$$x \in E_+(B, d) \implies x + jd \notin B \quad (1 \leq j \leq 2m), \\ x \in E_-(B, d) \implies x - jd \notin B \quad (1 \leq j \leq 2m).$$

If all the sets $E_+(B, d) + jd$ and $E_-(B, d) - jd$, $1 \leq j \leq m$, lie in a set U , then

$$2m b_B(d) \leq |U \setminus B|.$$

Proof. All the displayed translates avoid B by hypothesis. We prove they are pairwise disjoint. If $x + jd = y + kd$ with $x, y \in E_+$ and $j > k$, then $y = x + (j - k)d \in B$, contrary to non-return after the exit at x . The negative-sign case is identical. If $x + jd = y - kd$ with $x \in E_+$ and $y \in E_-$, then $y = x + (j + k)d \in B$, and $2 \leq j + k \leq 2m$, again a contradiction. Thus the $2m$ translates are disjoint. Each has cardinality $b_B(d)$ by Lemma II.128.3. Their union lies in $U \setminus B$. $\square$

The use of both signs is important. A one-sided argument loses a factor of two. The non-return statement is needed through $2m$, not merely through m , because a collision between opposite tubes uses $j + k$.

II.129.2 A boundary exit from a narrow Bohr strip cannot return quickly

Lemma II.129.2 (Coordinate non-return). *Let $B = B(\boldsymbol{\rho})$ with $0 < \rho_i \leq 1/4$, and suppose $\|\chi_i(d)\|_{\mathbb{T}} \leq \rho_i/(2m)$ for every i . Then the non-return hypotheses of Lemma II.129.1 hold.*

Proof. Take $x \in E_+(B, d)$ and choose a character witnessing $x + d \notin B$. Write $\rho = \rho_i$. There are real representatives

$$a \in [-\rho, \rho], \quad v \in [-\rho/(2m), \rho/(2m)]$$

for $\chi_i(x)$ and $\chi_i(d)$. Since $|a + v| \leq 3\rho/2 \leq 3/8 < 1/2$, the failure at $x + d$ means $|a + v| > \rho$ without any ambiguity from reduction modulo one.

Suppose first that $a + v > \rho$. Then $v > 0$, and for $1 \leq j \leq 2m$,

$$\rho < a + jv \leq \rho + 2m|v| \leq 2\rho \leq 1/2.$$

Therefore $\|\chi_i(x + jd)\|_{\mathbb{T}} > \rho$, so $x + jd \notin B$. The case $a + v < -\rho$ is the same with signs reversed. Applying the argument to $-d$ proves the assertion for E_- . The endpoint $2\rho = 1/2$ causes no problem, because $1/2 > \rho$. $\square$

Theorem II.129.3 (Intrinsic fixed-radius escape inequality). *Let G be a finite abelian group, $B = B(\boldsymbol{\chi}, \boldsymbol{\rho})$, and $0 < \rho_i \leq 1/4$. Suppose $m \geq 1$ is an integer and $v_i = \|\chi_i(d)\|_{\mathbb{T}} \leq \rho_i/(2m)$ for every i . Then*

$$2m b_B(d) \leq |B(\boldsymbol{\rho} + m\mathbf{v})| - |B| \leq |B(3\boldsymbol{\rho}/2)| - |B|. \quad (\text{II.129.1})$$

Proof. For $x \in B$ and $1 \leq j \leq m$, the triangle inequality on $\mathbb{T}$ gives

$$\|\chi_i(x \pm jd)\|_{\mathbb{T}} \leq \rho_i + jv_i \leq \rho_i + mv_i.$$

Consequently every tube in Lemma II.129.1 lies in $B(\boldsymbol{\rho} + m\mathbf{v})$. Lemma II.129.2 supplies non-return and hence the first inequality. The coordinatewise inequality $\rho_i + mv_i \leq 3\rho_i/2$ gives the second. $\square$

Remark II.129.4 (Degenerate steps and torsion). If d belongs to the common kernel of the characters, then $B + d = B$ and both exit sets are empty. No division by a step size is used in the proof. Similarly, no lower bound for the order of d is assumed: if an orbit is too short to support a nonempty exit tube, the non-return argument forces the relevant exit set to be empty. This is why the theorem does not need prime modulus or invertible frequencies.

II.130 Packing the enlarged Bohr set

Source 30, delivered Section 4.

Lemma II.130.1 (Three-cell packing). *If $0 < \rho_i \leq 1/4$, then*

$$|B(3\rho/2)| \leq 3^r |B(\rho)|. \quad (\text{II.130.1})$$

Proof. For each point of the enlarged Bohr set, choose the unique coordinate representative in $[-3\rho_i/2, 3\rho_i/2] \subset (-1/2, 1/2)$. Partition this interval into three disjoint intervals of diameter at most ρ_i ; assign the two internal endpoints consistently to one adjacent interval. The coordinate partitions give at most 3^r cells.

Fix a nonempty cell and a point y in it. If x is another point in the same cell, the chosen coordinate representatives differ by at most ρ_i . Their difference represents $\chi_i(x - y)$, so $x - y \in B$. The map $x \mapsto x - y$ is injective. Thus each cell has at most $|B|$ points, proving the bound. This argument counts points of G itself; it does not incorrectly discard multiplicities of the character map. $\square$

This is a finite box-packing argument of the same general kind as the standard Bohr doubling estimates; compare Green's Lemma 4.1 [124]. The constants here correspond specifically to enlargement by $3/2$ and cells of width ρ_i .

Theorem II.130.2 (Rank-explicit fixed-radius invariance). *Under the hypotheses of Theorem II.129.3,*

$$|B \cap (B + d)| \geq \left(1 - \frac{3^r - 1}{2m}\right) |B|. \quad (\text{II.130.2})$$

More precisely, with

$$D_B = \frac{|B(3\rho/2)|}{|B|},$$

one may replace $3^r - 1$ by $D_B - 1$.

Proof. Combine Theorem II.129.3 and Lemma II.130.1, then use $|B \cap (B + d)| = |B| - b_B(d)$. $\square$

If $D_B = 1$, every shift satisfying the step hypothesis in Theorem II.129.3 is an exact period of B . This conclusion requires no logarithm or division by $D_B - 1$.

Corollary II.130.3 (A floor-free relative-step estimate). *Let $t = \max_i \|\chi_i(d)\|_{\mathbb{T}} / \rho_i$. If $0 < t < 1/2$, then*

$$\frac{b_B(d)}{|B|} \leq \min \left\{ 1, \frac{D_B - 1}{2 \lfloor 1/(2t) \rfloor} \right\} \leq \min \left\{ 1, \frac{(D_B - 1)t}{1 - 2t} \right\}. \quad (\text{II.130.3})$$

In particular, the last numerator can be replaced by $(3^r - 1)t$. For $t = 0$, the defect is zero.

Proof. Set $m = \lfloor 1/(2t) \rfloor \geq 1$. Then $t \leq 1/(2m)$, so the main theorem applies. The inequality $m \geq 1/(2t) - 1$ yields $1/(2m) \leq t/(1 - 2t)$, and the trivial bound is $b_B(d) \leq |B|$. If $t = 0$, all the characters vanish on d and $B + d = B$. $\square$

The original estimate (F.35.1) has relative-step coefficient $2^{r+1}r\rho^{1-r}$. The new small-step coefficient is at most $3^r - 1$, independent of the width. The comparison is asserted on the new theorem's range $0 < \rho \leq 1/4$, not on all the larger range allowed by the old declaration. There is no hidden replacement of ρ by a regular radius.

II.130.1 Exact sharpness of the two constituent estimates

Proposition II.130.4 (The factor $2m$ is sharp). *For every positive even integer a and integer $N > 4a$, let $G = \mathbb{Z}/N\mathbb{Z}$, $\chi(x) = x/N$, $\rho = a/N$, $d = 1$, and $m = a/2$. Then equality holds in the first and second inequalities of (II.129.1).*

Proof. The set B is the interval $[-a, a]$ modulo N , of size $2a + 1$. Translation by one has defect one. Since a is even, the enlarged set is $[-3a/2, 3a/2]$, of size $3a + 1$. Thus

$$2mb_B(1) = a = (3a + 1) - (2a + 1).$$

Also $\rho + m/N = 3\rho/2$, so the direction-dependent enlargement is exactly the $3/2$ enlargement. $\square$

Proposition II.130.5 (The packing factor 3^r is sharp). *For every $r \geq 1$ there is a Bohr set with widths at most $1/4$ for which $|B(3\rho/2)| = 3^r|B|$.*

Proof. Take $G = (\mathbb{Z}/p\mathbb{Z})^r$ with $p \geq 5$, coordinate characters, and $\rho_i = 3/(4p)$. Then $B = \{0\}$, whereas the enlarged set consists of the coordinate choices $0, 1, -1$, because its width is $9/(8p)$. Its cardinality is 3^r . $\square$

The two sharpness assertions concern different inequalities. The packing example has no nonzero qualifying coordinate step. It does *not* show that the combined coefficient $3^r - 1$ in the translation estimate is optimal. Confusing these two forms of sharpness would erase the main remaining extremal problem.

II.131 Dense-domain extension of partial homomorphisms

Source 30, delivered Section 5.

Definition II.131.1 (Freiman homomorphism). A map $\psi : A \rightarrow H$ is a Freiman homomorphism of order k if every relation $a_1 + \cdots + a_k = b_1 + \cdots + b_k$ with entries in A , allowing repetition, implies $\psi(a_1) + \cdots + \psi(a_k) = \psi(b_1) + \cdots + \psi(b_k)$.

II.131.1 A general finite-pattern intersection bound

Lemma II.131.2 (Tree intersection budget). *Let $A \subseteq B \subseteq G$, $|B| > 0$, and let $v_1, \dots, v_q \in G$ be positions indexed by the vertices of a tree T . Repeated positions are permitted. Then*

$$\left| \bigcap_{j=1}^q (A + v_j) \right| \geq |B| - \sum_{\{i,j\} \in E(T)} b_B(v_i - v_j) - q|B \setminus A|. \quad (\text{II.131.1})$$

Proof. Root the tree and adjoin vertices so that every new vertex has a previously adjoined parent. The intersection initially has size $|B|$. Adjoining $B + v_j$ can remove at most

$$|(B + v_i) \setminus (B + v_j)| = b_B(v_i - v_j)$$

points, where i is the parent, since the current intersection is already a subset of $B + v_i$. Sum these losses. Finally, replacing each translate of B by the corresponding translate of A removes at most $|B \setminus A|$ points for each of the q labels. This remains valid when two positions coincide; then it merely overcounts a possible loss. $\square$

[write] 6 October 2026; sources 28–33. This is source 12’s Lemma II.34.3 (there with S_i in place of A and normalized losses b_B ; here b_B counts points).

II.131.2 Difference maps and higher-order relations

Theorem II.131.3 (Dense-domain extension with an explicit budget). *Let $A \subseteq B \subseteq G$, $|B| > 0$, and suppose*

$$|B \setminus A| \leq \eta|B|, \quad b_B(d) \leq \kappa|B| \quad (d \in C),$$

where $\eta, \kappa \geq 0$. Let $\psi : A \rightarrow H$ be a Freiman homomorphism of order two. For an integer $k \geq 2$, assume

$$(2k - 1)\kappa + 2k\eta < 1. \tag{II.131.2}$$

Then $C \subseteq A - A$, and there is a uniquely induced map $\psi_1 : C \rightarrow H$ satisfying

$$\psi_1(x - y) = \psi(x) - \psi(y) \quad (x, y \in A, x - y \in C). \tag{II.131.3}$$

This map is a Freiman homomorphism of order k .

Proof. For $d \in C$,

$$|A \cap (A + d)| \geq (1 - \kappa - 2\eta)|B| > 0, \tag{II.131.4}$$

because (II.131.2) implies $\kappa + 2\eta < 1$. Thus every d is a difference of two elements of A . If $x - y = x' - y'$, then $x + y' = x' + y$. The order-two hypothesis gives $\psi(x) - \psi(y) = \psi(x') - \psi(y')$. This proves existence and uniqueness of the induced difference map.

Now take $d_1, \dots, d_k, e_1, \dots, e_k \in C$ with equal total sums. Form two paths from zero to their common sum: their successive edge labels are $d_1, \dots, d_k$ and $e_1, \dots, e_k$. Regard the initial and final vertices as shared, but keep all other path positions as distinct formal labels even if their group values coincide. There are $2k$ vertex labels. Removing one edge from the resulting formal cycle gives a tree with $2k - 1$ edges. Each edge displacement is a member of C or its negative; Lemma II.128.3 gives the same defect for both signs.

Apply Lemma II.131.2 to the negatives of these positions. The intersection contains at least

$$[1 - (2k - 1)\kappa - 2k\eta]|B| > 0$$

points. Choose x in it. Then every vertex $x + v$ of both translated paths lies in A . Along every edge, (II.131.3) identifies its image increment with the corresponding value of ψ_1 . Telescoping along the two paths gives

$$\sum_{i=1}^k \psi_1(d_i) = \psi \left(x + \sum_i d_i \right) - \psi(x) = \sum_{i=1}^k \psi_1(e_i).$$

This is the required order- k relation. □

[write] 6 October 2026; sources 28–33. *Against source 12* (checked at the write). In source 12’s notation ε is the missing fraction η here and κ the normalized loss. Its Theorem II.35.2 gives the order-two conclusion under $3\eta + 2\kappa < 1$, sharp on its whole boundary (Theorem II.35.3); the budget here at $k = 2$ is $4\eta + 3\kappa < 1$, which implies it. For higher order compare source 12’s Theorem II.35.6. The budget here is a weaker form of source 12’s results, which source 30 does not cite.

For $k = 2$, the budget is $3\kappa + 4\eta < 1$. This is exactly the four-vertex mechanism in Gowers’s proof of Corollary 10.11, with the parameters left visible. We do not claim that this sufficient budget is necessary. For example, extra structure of A or of the map can make the extension valid far beyond what a union bound certifies.

II.132 Larger Bohr neighborhoods in Corollary 10.11

Source 30, delivered Section 6.

Corollary II.132.1 (Margin-preserving radius). *Let $r \geq 1$, $0 < \rho \leq 1/4$, and $B = B(\chi, \rho)$ in a finite abelian group. Suppose $A \subseteq B$, $|A| \geq 7|B|/8$, and $\psi : A \rightarrow H$ is a Freiman homomorphism of order two. Define*

$$\zeta_{\text{margin}} = \frac{\rho}{8(3^r - 1)}, \quad C_{\text{margin}} = B(\chi, \zeta_{\text{margin}}). \quad (\text{II.132.1})$$

Then $C_{\text{margin}} \subseteq A - A$ and the induced difference map is a Freiman homomorphism of order two on C_{margin} . Moreover, for every $d \in C_{\text{margin}}$,

$$|B \cap (B + d)| \geq \frac{7}{8}|B|, \quad |A \cap (A + d)| \geq \frac{5}{8}|B|.$$

For each four-term relation on C_{margin} , the four-vertex construction has at least $|B|/8$ possible base points in A .

Proof. Set $m = 4(3^r - 1)$, which is a positive integer. The shift radius in (II.132.1) is exactly $\rho/(2m)$. Theorem II.130.2 gives $\kappa = 1/8$. Use Theorem II.131.3 with $k = 2$ and $\eta = 1/8$; its budget is $3/8 + 4/8 = 7/8 < 1$. The pair and four-vertex lower bounds follow from the same counting argument. $\square$

Corollary II.132.2 (A larger extension-only radius). *Under the hypotheses of Corollary II.132.1, the extension conclusion also holds on*

$$C_{\text{ext}} = B(\chi, \zeta_{\text{ext}}), \quad \zeta_{\text{ext}} = \frac{\rho}{6 \cdot 3^r - 4}. \quad (\text{II.132.2})$$

For $c_r = 3^r - 1$, the four-vertex base-point count is at least $|B|/(6c_r + 2)$, and the pair overlap is greater than $7|B|/12$.

Proof. Choose $m = 3c_r + 1$. Then

$$\kappa = \frac{c_r}{6c_r + 2} < \frac{1}{6}, \quad 3\kappa + 4\eta = \frac{3c_r}{6c_r + 2} + \frac{1}{2} < 1.$$

The residual four-vertex fraction is $1 - 3\kappa - 1/2 = 1/(6c_r + 2)$. The pair fraction is $3/4 - \kappa > 7/12$. $\square$

[write] **6 October 2026; sources 28–33.** *A larger fixed radius, by composition with source 12 (proved at the write). Theorem II.130.2 gives normalized losses $b_B(d)/|B| \leq c_r/(2m)$ on the symmetric set $B(\chi, \rho/(2m)) \ni 0$. Source 12's Theorem II.35.2 with $\varepsilon = \eta = 1/8$ needs $2\kappa < 5/8$, that is $m > 8c_r/5$. So the extension conclusion of this corollary holds on $B(\chi, \rho/(2m))$ with $m = \lfloor 8c_r/5 \rfloor + 1$, about $5\rho/(16c_r)$, against $\rho/(6 \cdot 3^r - 4) = \rho/(2(3c_r + 1))$ here: larger by the factor $(3c_r + 1)/m$, which is 1.75 at $r = 1$, 1.92 at $r = 2$ and tends to 1.875 (exact for $r \leq 10$). The radius is still fixed; the four-vertex base-point count is then $(1 - 3/8 - 2\kappa)|B| > 0$.*

The two radii serve different interfaces. The margin-preserving radius can be used wherever the downstream proof explicitly consumes the $5/8$ pair overlap. The extension-only radius requires recalculation of that downstream budget; it must not be substituted while silently retaining $5/8$.

II.132.1 Quantitative gain over the corrected baseline

Combining (F.35.2) and (II.132.1) gives

$$\frac{\zeta_{\text{margin}}}{\zeta_{\text{old}}} = \frac{2^{r+1}r}{3^r - 1} \rho^{1-r}. \quad (\text{II.132.3})$$

This ratio is at least two for every $r \geq 1$ and $0 < \rho \leq 1/4$. Indeed, $c_1 = 2$ and $c_{r+1} = 3c_r + 2 \leq 4c_r$, so $3^r - 1 \leq 2 \cdot 4^{r-1}$. Together with $\rho^{1-r} \geq 4^{r-1}$, this gives

$$\frac{\zeta_{\text{margin}}}{\zeta_{\text{old}}} \geq 2^r r \geq 2.$$

The equality at $r = 1$ is the radius comparison, not an assertion that the radius itself is optimal. For $r > 1$, the gain diverges as $\rho \rightarrow 0$.

r	ζ_{old} at $\rho = 10^{-2}$	ζ_{margin}	gain
1	$3.125 \cdot 10^{-4}$	$6.250 \cdot 10^{-4}$	2
2	$7.813 \cdot 10^{-7}$	$1.563 \cdot 10^{-4}$	200
3	$2.604 \cdot 10^{-9}$	$4.808 \cdot 10^{-5}$	$1.846 \cdot 10^4$
5	$3.906 \cdot 10^{-14}$	$5.165 \cdot 10^{-6}$	$1.322 \cdot 10^8$
10	$6.104 \cdot 10^{-26}$	$2.117 \cdot 10^{-8}$	$3.468 \cdot 10^{17}$

These are rounded displays of rational numbers. The exact fractions and comparison identities are included in the verification output. A large radius ratio does not by itself imply a comparable improvement in a final Szemerédi bound.

II.132.2 Higher order and measured local growth

The same calculation gives a useful family of choices. Suppose $|A| \geq (1 - \eta)|B|$, $\eta < 1/(2k)$, and $k \geq 2$. Any integer

$$m > \frac{(2k - 1)(D_B - 1)}{2(1 - 2k\eta)} \quad (\text{II.132.4})$$

makes the induced difference map an order- k Freiman homomorphism on $B(\rho/(2m))$. One may replace $D_B - 1$ by $3^r - 1$. For example, density $7/8$ and $k = 3$ are covered by $m = 10(3^r - 1) + 1$. This illustrates a gain that is qualitatively different from merely changing a constant: the same order-two map on a very dense domain acquires higher-order relations on a sufficiently small neighborhood.

For measured local growth, a margin-preserving choice is $m = \max\{1, \lceil 4(D_B - 1) \rceil\}$. Its radius is $\rho/(2m)$; in the rectangular case all widths are divided by $2m$. The ceiling and the maximum are needed to handle integral and degenerate cases correctly.

II.133 A density-sensitive compatibility step for Lemma 10.12

Source 30, delivered Section 7.

We now isolate the finite-fibre argument that transfers an induced map back to the multifunction domain. Keeping the actual size of A in the exception count improves a local numerical margin. This argument is independent of how the map ψ was obtained.

Let X be a finite set with an index map $r_X : X \rightarrow G$, write $X_g = \{z \in X : r_X(z) = g\}$, and let $\phi : X \rightarrow H$. Let $W_1 \subseteq X$, $A \subseteq B$, and $\psi : A \rightarrow H$. Assume that for each $u \in W_1$ there is a set $A_u \subseteq A$ with

$$|A_u| \geq (1 - \sigma)|A|, \quad (\text{II.133.1})$$

and that, for each $a \in A_u$, the fibre $X_{r_X(u)+a}$ is nonempty and at least a proportion $1 - \theta$ of it satisfies

$$\phi(z) - \phi(u) = \psi(a). \quad (\text{II.133.2})$$

Theorem II.133.1 (Compatibility with exact direction-density budget). *Let $a_0 = |A|/|B|$, $0 \leq \theta < 1/2$, and $\sigma \geq 0$. Suppose (II.133.1)–(II.133.2) hold, and $\psi_1 : C \rightarrow H$ is induced by ψ as in (II.131.3). If $b_B(c) \leq \kappa|B|$ for $c \in C$ and*

$$2(1 - \sigma)a_0 > 1 + \kappa, \quad (\text{II.133.3})$$

then, for $u, v \in W_1$ with $c = r_X(u) - r_X(v) \in C$,

$$\phi(u) - \phi(v) = \psi_1(c).$$

Proof. The set of directions a with $a \in A$ and $a + c \in A$ has cardinality at least

$$|B \cap (B - c)| - 2|B \setminus A| \geq 2|A| - (1 + \kappa)|B|.$$

Requiring additionally $a \in A_u$ and $a + c \in A_v$ removes at most $2\sigma|A|$ directions. The resulting count is positive by (II.133.3); choose such an a . The two target fibres coincide because $r_X(u) + a = r_X(v) + (a + c)$, and they are nonempty by hypothesis. The two good-value sets inside this common fibre each have relative size at least $1 - \theta$. Since $2\theta < 1$, their intersection is nonempty. Choose z in the intersection. Then

$$\phi(z) - \phi(u) = \psi(a), \quad \phi(z) - \phi(v) = \psi(a + c).$$

Subtract and use the induced difference identity to obtain the result. $\square$

[write] **6 October 2026; sources 28–33.** The condition (II.133.3) is that of source 12’s Lemma II.39.1, $2(1 - q)a > 1 + \kappa$ with $q = \sigma$ and $a = a_0$, proved by the same intersection. The 5/14 of the corollary below improves the formal lemma_10_12’s $\sqrt{\theta} < 5/16$ (Section10.lean, line 401, at HEAD 5d4718772), as the text says.

Corollary II.133.2 (The 5/14 margin). *In the setting of Theorem II.133.1, suppose $|A| \geq 7|B|/8$, $\kappa \leq 1/8$, and $\sigma = \sqrt{\theta}$ with $0 \leq \theta$ and $\sqrt{\theta} < 5/14$. Then the compatibility conclusion holds. In particular it applies on C_{margin} from Corollary II.132.1 whenever the stated fibre hypotheses hold.*

Proof. The condition $\sqrt{\theta} < 5/14$ implies $\theta < 25/196 < 1/2$ and

$$2(1 - \sqrt{\theta})\frac{7}{8} > 2\left(1 - \frac{5}{14}\right)\frac{7}{8} = \frac{9}{8} \geq 1 + \kappa.$$

Apply Theorem II.133.1. $\square$

The repository's corrected Lemma 10.12 explicitly uses $\sqrt{\theta} < 5/16$, together with nonempty target fibres [138]. That budget comes from a $5|B|/8$ pair overlap and two exception counts each bounded by $\sqrt{\theta}|B|$. The refined count above retains $|A|$ until the end, giving $5/14$ instead. This does not assert that $5/14$ is an optimal compatibility threshold. Nor does it remove the nonempty-fibre condition: on an empty fibre a relative “almost every” statement can be vacuous.

For the larger extension-only radius, the bound $\kappa < 1/6$ gives the convenient sufficient condition $\sqrt{\theta} \leq 1/3$ under the same nonempty-fibre assumptions. Indeed, $2(1 - 1/3)(7/8) = 7/6 > 1 + \kappa$. Thus one may trade some compatibility margin for a larger neighborhood, with the tradeoff recorded explicitly.

II.134 Further consequences without radius regularization

Source 30, delivered Section 8.

II.134.1 Uniform measures and bounded observables

Let $\mu_B = \mathbf{1}_B/|B|$ be the uniform probability mass function on B , and let $\tau_d\mu_B(x) = \mu_B(x - d)$. With total variation defined as one half of the counting ℓ^1 distance,

$$\mathrm{TV}(\mu_B, \tau_d\mu_B) = \frac{b_B(d)}{|B|}. \quad (\text{II.134.1})$$

This is immediate from Lemma II.128.3. Therefore the main theorem gives

$$\|\tau_d\mu_B - \mu_B\|_1 \leq \frac{D_B - 1}{m}.$$

For an observable $0 \leq f \leq 1$, the difference between its expectations under the two measures is at most $(D_B - 1)/(2m)$. To see this without invoking a general duality theorem, separate the positive and negative parts of the signed measure: both have mass $b_B(d)/|B|$. For complex f with $|f| \leq 1$, the triangle inequality gives twice that bound.

If convolution is defined by $(f * \mu)(x) = \sum_y f(x - y)\mu(y)$, then

$$\|f * (\tau_d\mu_B - \mu_B)\|_\infty \leq \|f\|_\infty \frac{D_B - 1}{m}.$$

No smoothed cutoff is introduced. These estimates can be useful when a proof must retain exactly the original indicator of a Bohr set.

II.134.2 Long orbit segments inside the original set

For an integer $L \geq 0$, define

$$B_{L,d} = \{x \in G : x + jd \in B \text{ for } -L \leq j \leq L\}.$$

Apply Lemma II.131.2 to the path of $2L + 1$ positions, with $A = B$:

$$|B_{L,d}| \geq |B| - 2L b_B(d) \geq \left(1 - \frac{L(D_B - 1)}{m}\right) |B|. \quad (\text{II.134.2})$$

This counts base points of orbit segments. It does not claim the segment has $2L + 1$ distinct points; that requires a separate lower bound on the order of d . The statement remains valid when the orbit repeats.

More generally, any finite pattern equipped with a tree of small character-displacement edges has a large intersection of its translated Bohr sets, with the sum of the intrinsic edge defects controlling the loss. This is an interface for higher-order local identities; the tree budget should be carried to the final application rather than replaced prematurely by a worst-case cardinality factor.

II.135 Exponential-rank obstructions, including cyclic groups

Source 30, delivered Section 9.

The factor 3^r suggests asking whether a more careful argument could give a polynomial in r instead. We show that no universal polynomial-rank translation bound is possible for arbitrary fixed input Bohr sets, even when the ambient group is required to be cyclic. The meaning of “fixed” here is that the radius is not adjusted within an instance; the group and radius may vary across a family of counterexamples.

II.135.1 A simple product-group model

Proposition II.135.1 (A line with a transverse spike). *Fix $p \geq 5$, $s \geq 1$, and $L \geq 2$, and put $q = Lp - 1$. On $G = \mathbb{Z}/q\mathbb{Z} \times (\mathbb{Z}/p\mathbb{Z})^s$, use the $2s + 1$ characters*

$$\chi_0(x, \mathbf{y}) = x/q, \quad \chi_i^\pm(x, \mathbf{y}) = y_i/p \pm x/q \quad (1 \leq i \leq s)$$

and common width $\rho = 1/p$. Then

$$B = (I \times \{\mathbf{0}\}) \cup (\{0\} \times \{-1, 0, 1\}^s), \quad I = [-(L-1), L-1] \pmod{q}. \quad (\text{II.135.1})$$

For $d = (1, \mathbf{0})$,

$$|B| = 2L - 2 + 3^s, \quad b_B(d) = 3^s, \quad \max_x \frac{\|\chi(d)\|_{\mathbb{T}}}{\rho} = \frac{p}{q}. \quad (\text{II.135.2})$$

Consequently,

$$\frac{b_B(d)/|B|}{p/q} \longrightarrow \frac{3^s}{2} \quad \text{as } L \rightarrow \infty.$$

Proof. Choose a representative $t = x/q$ with $|t| \leq \rho$, as forced by χ_0 . If both $\|y_i/p + t\|_{\mathbb{T}}$ and $\|y_i/p - t\|_{\mathbb{T}}$ are at most ρ , there are real lifts $u, v \in [-\rho, \rho]$ of these two elements. Now $u - v$ represents $2t$. Since $|u - v - 2t| \leq 4\rho < 1$, the integral difference must be zero. Thus $u = z + t$ and $v = z - t$ for a lift z of y_i/p , and

$$|z| + |t| = \max\{|z + t|, |z - t|\} \leq \rho.$$

In particular $z \in \{0, 1/p, -1/p\}$. If $z \neq 0$, the inequality forces $t = 0$. If all z vanish, the only restriction is $|x| \leq \lfloor q/p \rfloor = L - 1$. This proves (II.135.1). The nonzero transverse fibres of B have only the point $x = 0$; their translates have only $x = 1$. The axial fibre is an interval whose overlap with its translate loses one point. Hence the defect is one per transverse choice, namely 3^s , and the size is as stated. Finally, $q/p = L - 1/p$, so the ratio tends to $3^s/2$. $\square$

The finite products with $s > 1$ are not cyclic. The following construction removes that limitation rather than silently identifying a product with a cyclic group when its factors are not coprime.

II.135.2 An explicit Chinese-remainder realization

For $s \geq 1$ and integer $L \geq 9$, define

$$D = s!, \quad R_i = LiD + 1, \quad P = \left(\prod_{i=1}^s R_i \right)^2, \quad P_i = P + iD, \quad q = LP - 1. \quad (\text{II.135.3})$$

Lemma II.135.2 (Arithmetic of the parameters). *The integers $q, P_1, \dots, P_s$ are pairwise coprime. Also*

$$P \geq 100, \quad P > LsD, \quad 0 < \frac{1}{P} - \frac{1}{P_i} < \frac{1}{q}.$$

Proof. Each R_i is congruent to one modulo D , so $P \equiv 1 \pmod{D}$. If a prime divides both P_i and P_j for $i < j$, it divides $(j - i)D$. Every prime divisor of this number divides D : primes dividing $j - i \leq s - 1$ are among the primes dividing $s!$. But $P_i \equiv 1 \pmod{D}$, a contradiction. For $s = 1$ this pairwise argument is vacuous.

A common divisor of q and P_i divides $LP_i - q = LiD + 1 = R_i$. Since $R_i \mid P$,

$$\gcd(P_i, R_i) = \gcd(iD, R_i) = 1.$$

Thus $\gcd(q, P_i) = 1$. The product defining P contains $R_s = LsD + 1$, and its square is larger still, so $P > LsD$. The smallest case gives $P \geq (L + 1)^2 \geq 100$. Finally,

$$0 < \frac{1}{P} - \frac{1}{P_i} = \frac{iD}{PP_i} \leq \frac{sD}{P^2} < \frac{1}{LP} < \frac{1}{q}.$$

□

Theorem II.135.3 (Cyclic spiked Bohr sets). *Use the group*

$$G = \mathbb{Z}/q\mathbb{Z} \times \prod_{i=1}^s \mathbb{Z}/P_i\mathbb{Z}$$

with the parameters (II.135.3). It is cyclic of order $N = q \prod_i P_i$. With characters

$$\chi_0(x, \mathbf{y}) = x/q, \quad \chi_i^\pm(x, \mathbf{y}) = y_i/P_i \pm x/q$$

and common width $\rho = 1/P$, its Bohr set is exactly (II.135.1), with $\{-1, 0, 1\}$ now interpreted in each $\mathbb{Z}/P_i\mathbb{Z}$. For $d = (1, \mathbf{0})$,

$$|B| = 2L - 2 + 3^s, \quad b_B(d) = 3^s, \quad t = \max_{\chi} \frac{\|\chi(d)\|_{\mathbb{T}}}{\rho} = \frac{P}{q}, \quad (\text{II.135.4})$$

and

$$\frac{b_B(d)/|B|}{t} = \frac{3^s(L - 1/P)}{2L - 2 + 3^s} \longrightarrow \frac{3^s}{2} \quad (L \rightarrow \infty). \quad (\text{II.135.5})$$

Proof. Cyclicity follows from Lemma II.135.2 and the Chinese remainder theorem. The base character gives a lift $t_x = x/q$ with $|t_x| \leq 1/P$. As in Proposition II.135.1, the two constraints for a fixed i imply the existence of a lift z_i of y_i/P_i such that

$$|z_i| + |t_x| \leq 1/P. \quad (\text{II.135.6})$$

The lift argument applies since $4/P < 1$. As $P_i/P = 1 + iD/P < 2$, the only possible values are $z_i = 0, \pm 1/P_i$. If any z_i is nonzero, then

$$|t_x| \leq 1/P - 1/P_i < 1/q.$$

But t_x is an integer multiple of $1/q$, so $x = 0$. Conversely, at $x = 0$ all 3^s transverse choices are allowed because $P_i > P$. When all $z_i = 0$, the remaining restriction is $|x| \leq \lfloor q/P \rfloor = L - 1$. This proves the exact description. The cardinality and defect follow by the same interval calculation as in Proposition II.135.1. All character values at d have norm $1/q$, giving $t = P/q$. Since $P \geq 100$, the limit follows directly from the displayed rational formula. $\square$

The construction can be written entirely in $\mathbb{Z}/N\mathbb{Z}$. Under the map $z \mapsto (z \bmod q, z \bmod P_1, \dots, z \bmod P_s)$, the frequencies are

$$K = \left\{ \frac{N}{q} \right\} \cup \left\{ \frac{N}{P_i} + \frac{N}{q}, \frac{N}{P_i} - \frac{N}{q} : 1 \leq i \leq s \right\} \pmod{N}. \quad (\text{II.135.7})$$

They are distinct: restrictions to the individual direct factors separate different i , and restriction to the q -factor separates the two signs, since $q > 2$. The shift is the unique residue with $d \equiv 1 \pmod{q}$ and $d \equiv 0 \pmod{P_i}$ for every i . Thus no abstract existence of unspecified frequencies is needed.

II.135.3 The precise obstruction to polynomial rank

Definition II.135.4 (Small-step extremal coefficient). Let $S_r(\varepsilon)$ be the supremum of

$$\frac{b_B(d)}{t|B|}, \quad t = \max_i \frac{\|\chi_i(d)\|_{\mathbb{T}}}{\rho_i},$$

over finite cyclic groups, Bohr sets with r listed characters and $0 < \rho_i \leq 1/4$, and shifts with $0 < t \leq \varepsilon$. Set $\mathcal{L}_r = \lim_{\varepsilon \downarrow 0} S_r(\varepsilon)$. The limit exists as an extended real number by monotonicity.

Corollary II.135.5 (Exponential dependence is necessary). *For every $r \geq 1$, $\mathcal{L}_r \leq 3^r - 1$. For $s \geq 1$,*

$$\frac{1}{2} 3^s \leq \mathcal{L}_{2s+1} \leq 3^{2s+1} - 1. \quad (\text{II.135.8})$$

Consequently no coefficient bounded by a polynomial in r can give a universal linear relative-shift bound, even restricted to cyclic groups and sufficiently small shifts.

Proof. Corollary II.130.3 gives $S_r(\varepsilon) \leq (3^r - 1)/(1 - 2\varepsilon)$ for $0 < \varepsilon < 1/2$. Take the limit. For the lower bound, the cyclic family in Theorem II.135.3 has $t = P/(LP - 1) \rightarrow 0$ and its normalized ratio tends to $3^s/2$. Thus every sufficiently small-step supremum is at least this limiting value. Exponential growth on the odd subsequence already rules out any polynomial bound valid for all r . $\square$

These examples do not settle the prime-order case. They also do not settle the optimal exponential base: the lower bound grows like $(\sqrt{3})^r$, whereas the upper bound grows like 3^r . The local-growth formulation of Theorem II.129.3 remains preferable when more structure of the individual Bohr set is available.

II.136 Source 33's chapter: high-agreement repair of respected arrangements [write]

Added in the sixth write, 6 October 2026.

The next seven sections are source 33's Sections 2–8, on Gowers's arrangements of Section 12 when almost all of them are respected on a full product. *Logically the chapter continues source 14's equality theory for arrangement moments* (Section II.44) *and source 15's chapter* (Section II.45); it is printed here because no number printed earlier may move. Source 33's introduction, with its main theorems, is Section F.38; its questions and integration notes are in Part V.

What is proved, in one paragraph. The exact models of the arrangement test are $a(x) + B(x, y) + \ell(y)$ with an arbitrary row function. If the failure probability of the test with s vertical differences is $\varepsilon_s \leq 1/(240(s-1))$, the distance to the models is at most $R_s(\varepsilon_s) = \varepsilon_s/s + 2(s-1)\varepsilon_s^2/s^2 + O_s(\varepsilon_s^3)$, both coefficients optimal over finite abelian groups and the first already over $\mathbb{F}_p$ (Theorem F.38.2, with Sections II.140–II.141); for $s = 4$ the whole cubic $4\delta - 24\delta^2 + 32\delta^3$ is sharp (Theorem F.38.3). Along the way: a cycle lemma giving $\max\{\theta^7, 4\theta - 3\}N^{32}$ respected eight-arrangements in Gowers's Lemma 12.3 (Corollary II.137.6), a coarse decoder (Theorem II.139.1), erasures (Theorem II.142.1) and an exact cyclic decoder. The majority correction is credited to Blum, Luby and Rubinfeld. The sparse case, the one Gowers needs, is open (item II.25).

Relation to the formal statements (checked at HEAD 5d4718772). `lemma_12_3` (Sections12_13.lean, line 151) has the exact companion `lemma_12_3_holds` (Proofs12HigherArrangements.lean, line 334); the linear branch is a formalization candidate subject to matching the counts' conventions (note after the corollary). `lemma_12_5` and `lemma_12_6` (lines 170, 183) have no counterpart of the repair theorem. Nothing in this chapter is formalized.

II.137 Sampling, independence, and Fourier normalization

Source 33, delivered Section 2.

II.137.1 Uniform sampling on a balanced relation

Let $\mathcal{X}_s \subset G^s$ be the solutions of (F.38.1). One may sample them by choosing the first $s-1$ coordinates independently and solving uniquely for the last. Thus $|\mathcal{X}_s| = |G|^{s-1}$.

Lemma II.137.1 (Proper-coordinate independence). *For every proper subset $I \subsetneq \{1, \dots, s\}$, the coordinates $(x_i)_{i \in I}$ of a uniform element of $\mathcal{X}_s$ are independent and uniform. Consequently, conditional on any fixed h , the pairs (x_i, y_i) in any proper subset of positions are independent and uniform on $G \times H$.*

Proof. Fix the coordinates in I . Among the remaining positions choose one, j , to solve for. The other $s - |I| - 1$ coordinates may be chosen freely. The coefficient of x_j in the balanced relation is 1 or -1 , both invertible as endomorphisms of any abelian group. Hence there are exactly $|G|^{s-|I|-1}$ completions, independently of the specified coordinates. The statement about pairs follows because the y_i are independent of all x_i and of one another. $\square$

The total number of ordered arrangements is

$$|G|^{s-1}|H|^{s+1}. \tag{II.137.1}$$

In particular, when $G = H = \mathbb{Z}_N$, it is N^{2s} : N^8 for $s = 4$ and N^{32} for $s = 16$. The s pairs are not generally jointly independent; the missing last degree of independence will be important in the sharp local estimate.

II.137.2 An exact moment identity

For a finite abelian group L , write $\widehat{L}$ for its character group, with characters taking values on the complex unit circle. On G use

$$\widehat{g}(\xi) = \mathbb{E}_{x \in G} g(x) \overline{\xi(x)}.$$

Thus Fourier coefficients are normalized, while sums over frequencies below are not. For $\chi \in \widehat{K}$ and $h \in H$, define

$$g_{\chi,h}(x) = \mathbb{E}_{y \in H} \chi(D_h \phi(x, y)). \quad (\text{II.137.2})$$

Proposition II.137.2 (Moment formula). *For every even $s \geq 4$,*

$$1 - \varepsilon_s(\phi) = \mathbb{E}_{h \in H} \mathbb{E}_{\chi \in \widehat{K}} \sum_{\xi \in \widehat{G}} |\widehat{g}_{\chi,h}(\xi)|^s. \quad (\text{II.137.3})$$

Consequently $\varepsilon_{s'}(\phi) \geq \varepsilon_s(\phi)$ for even $s' \geq s \geq 4$.

Proof. Character orthogonality on K writes the indicator that the signed sum in (F.38.2) is zero as its character average. For fixed h and χ , averaging the independent y_i leaves

$$\mathbb{E}_{\mathbf{x} \in \mathcal{X}_s} \prod_{i \leq s/2} g_{\chi,h}(x_i) \prod_{i > s/2} \overline{g_{\chi,h}(x_i)}.$$

Insert the Fourier expansion of each factor. Orthogonality in the $s - 1$ free horizontal coordinates forces the same Fourier frequency in all factors. The resulting sum is $\sum_{\xi} |\widehat{g}_{\chi,h}(\xi)|^s$. This proves (II.137.3). Since $|\widehat{g}_{\chi,h}(\xi)| \leq \mathbb{E}_x |g_{\chi,h}(x)| \leq 1$, each term decreases when the even exponent increases. $\square$

Remark II.137.3. The use of a normalized character average over $\widehat{K}$ but an unnormalized sum over $\widehat{G}$ is deliberate. The constant function has just one nonzero Fourier coefficient, equal to one, so the right side of (II.137.3) is one for a perfectly respected test.

II.137.3 A cycle bound and a direct refinement of Lemma 12.3

The usual moment interpolation can be supplemented by information that each individual Fourier square-mass is at most one. This yields a useful linear bound near complete agreement.

Lemma II.137.4 (Cycle deficit inequality). *If $a_1, \dots, a_t \geq 0$, $\sum_i a_i \leq 1$, and $r \geq 2$ is an integer, then*

$$1 - \sum_i a_i^r \leq \frac{r}{2} \left(1 - \sum_i a_i^2 \right). \quad (\text{II.137.4})$$

Proof. First suppose the a_i sum to one. Draw r independent labels $X_1, \dots, X_r$ with this distribution and set $X_{r+1} = X_1$. Unless all labels agree, at least two neighboring pairs on this cycle disagree. This also holds for $r = 2$, counting both directed edges. Therefore

$$2\mathbb{P}[\text{not all } X_i \text{ agree}] \leq \sum_{j=1}^r \mathbb{P}[X_j \neq X_{j+1}],$$

which is (II.137.4). For total mass $A \leq 1$, adjoin one label of mass $b = 1 - A$. Applying the probability-vector case and rearranging gives

$$1 - \sum_i a_i^r \leq \frac{r}{2} \left(1 - \sum_i a_i^2 \right) + b^r - \frac{r}{2} b^2.$$

Since $0 \leq b \leq 1$ and $r \geq 2$, the last two terms sum to a nonpositive quantity. This proves the subprobability version without a limit. $\square$

Proposition II.137.5 (Sharp cross-order defect factor). *For every full-domain function and every even $s \geq 4$,*

$$\varepsilon_4(\phi) \leq \varepsilon_s(\phi) \leq \frac{s}{4} \varepsilon_4(\phi). \quad (\text{II.137.5})$$

The upper coefficient $s/4$ is optimal, even on common odd prime cyclic groups. It requires no small-error hypothesis.

Proof. The first inequality is Proposition II.137.2. For fixed h, χ , set $a_\xi = |\widehat{g}_{\chi,h}(\xi)|^2$. Parseval gives $\sum_\xi a_\xi = \mathbb{E}_x |g_{\chi,h}(x)|^2 \leq 1$. Apply Lemma II.137.4 with $r = s/2$, then average over h, χ and use (II.137.3). Optimality follows from either extremal family of Section II.141, where $\varepsilon_s/\varepsilon_4 \rightarrow s/4$. $\square$

Corollary II.137.6 (Supplement to Gowers's arrangement amplification). *Let $A \subset \mathbb{Z}_N^2$, and let $\phi : A \rightarrow \mathbb{Z}_N$. Suppose there are θN^8 respected pairs of vertical parallelograms, counted with Gowers's order and repetition conventions. Then there are at least*

$$\max\{\theta^7, 4\theta - 3\} N^{32} \quad (\text{II.137.6})$$

respected eight-arrangements. More generally, if θ_s denotes the number of respected $s/2$ -arrangements divided by N^{2s} , then

$$\theta_s \geq \max \left\{ \theta_4^{s/2-1}, \frac{s}{4} \theta_4 - \left(\frac{s}{4} - 1 \right) \right\}. \quad (\text{II.137.7})$$

The θ_s are ambient-normalized counts, not conditional fractions among arrangements contained in A .

Proof. In (II.137.2), replace the summand by zero unless both (x, y) and $(x, y + h)$ lie in A . Orthogonality still gives the moment formula for θ_s , and $|g_{\chi,h}(x)| \leq 1$ still holds. The cycle argument proves the second term of (II.137.7). For completeness, let $a_\xi = |\widehat{g}_{\chi,h}(\xi)|^2$ and $M_j = \mathbb{E}_{h,\chi} \sum_\xi a_\xi^j$. Then $M_1 \leq 1$, $M_2 = \theta_4$, and $M_{s/2} = \theta_s$. Hölder interpolation gives, for $r = s/2 > 2$,

$$M_2 \leq M_1^{(r-2)/(r-1)} M_r^{1/(r-1)}.$$

If $M_1 = 0$ everything is zero; otherwise $M_r \geq M_2^{r-1}/M_1^{r-2} \geq M_2^{-1}$. For $r = 2$ the statement is an identity. Taking $s = 16$ gives (II.137.6). $\square$

[write] **6 October 2026; sources 28–33.** *Against the formal statement* (checked at the write at HEAD 5d4718772). `lemma_12_3` (Sections12_13.lean, line 151) asks $\theta^7 N^{32} \leq \text{respectedArrangementCount}$ 8 from $\theta N^8 \leq \text{parallelogramPairCount}$; its exact companion is `lemma_12_3_holds` (Proofs12HigherArrangements.lean, line 334). The linear branch $4\theta - 3$ exceeds θ^7 exactly for $\theta \in (0.8045539\dots, 1)$. That this corollary strengthens the formal conclusion under the same hypothesis depends on matching the repetition and $h = 0$ conventions of the two formal counts with those here, which is not checked: a formalization candidate, not formalized. The branch θ^7 is Gowers's; source 15's Theorem II.47.3 is its collision-sensitive form.

The first term in (II.137.6) is Gowers's Lemma 12.3. The second improves its near-one regime without changing its hypotheses; for instance, at $\theta = 9/10$ it gives $3/5$ in place of $(9/10)^7 = 4782969/10^7$. The linear coefficient four is sharp at $\theta = 1$, by Proposition II.137.5. For smaller θ the original power bound can be stronger, which is why the maximum is retained. This elementary cycle estimate is supplementary; we do not claim it is the optimal complete moment envelope.

II.137.4 Exact convolution counts

For cyclic groups $G = \mathbb{Z}_n$, $H = \mathbb{Z}_m$, $K = \mathbb{Z}_q$, put

$$M_h(x, z) = |\{y \in \mathbb{Z}_m : D_h \phi(x, y) = z\}|.$$

Let convolution on $\mathbb{Z}_n \times \mathbb{Z}_q$ be unnormalized. Then

$$\text{Res}_s(\phi) = \sum_{h \in \mathbb{Z}_m} \sum_{(x, z) \in \mathbb{Z}_n \times \mathbb{Z}_q} (M_h^{*(s/2)}(x, z))^2 \quad (\text{II.137.8})$$

is the exact number of respected arrangements. Indeed, the convolution coefficient counts ordered half-arrangements with prescribed horizontal sum and prescribed derivative sum; squaring matches two halves with the same two sums. Dividing by $n^{s-1}m^{s+1}$ gives $1 - \varepsilon_s$. This identity permits exact tests with large s without enumerating all $2s$ point occurrences and without rounding Fourier coefficients.

II.138 A noisy homomorphism lemma with auxiliary labels

Source 33, delivered Section 3.

The first objective is only a coarse linear repair estimate. The sharp coefficient will come later from the geometry of a nearest model. A small technical point matters: each query in an arrangement has an independent vertical base point. Randomly choosing one representative value for each horizontal coordinate would not preserve this query law when coordinates repeat. The next lemma keeps those auxiliary labels independent throughout.

Lemma II.138.1 (Noisy homomorphism correction). *Let U, V be finite abelian groups, let S be a finite nonempty set, and let $F : U \times S \rightarrow V$. Suppose*

$$e = \mathbb{P}_{u,v,\alpha,\beta,\gamma}[F(u, \alpha) + F(v, \beta) \neq F(u + v, \gamma)] < \frac{1}{6}. \quad (\text{II.138.1})$$

Then there is a homomorphism $L : U \rightarrow V$ such that

$$\mathbb{P}_{u,\alpha}[F(u, \alpha) \neq L(u)] \leq \frac{e}{1 - 2e}. \quad (\text{II.138.2})$$

Proof. For a fixed $z \in U$, consider the random variable

$$A_z(y; \alpha, \beta) = F(z + y, \alpha) - F(y, \beta).$$

Two independent copies of A_z agree with probability at least $1 - 2e$. To prove this, introduce an independent label γ and compare

$$\begin{aligned} F(z + y, \alpha) + F(y', \beta') &= F(z + y + y', \gamma), \\ F(y, \beta) + F(z + y', \alpha') &= F(z + y + y', \gamma). \end{aligned}$$

Each equality has failure probability e , since its two group inputs are independent uniform elements even with z fixed. If both hold, $A_z(y; \alpha, \beta) = A_z(y'; \alpha', \beta')$.

Write $p_z(t) = \mathbb{P}[A_z = t]$. Then

$$\max_t p_z(t) \geq \sum_t p_z(t)^2 \geq 1 - 2e.$$

In particular, A_z has a unique most likely value $L(z)$, and

$$\mathbb{P}[A_z = L(z)] \geq 1 - 2e \quad \text{for every } z. \quad (\text{II.138.3})$$

There is a pointwise identity

$$A_{z+z'}(y; \alpha, \beta) = A_z(y + z'; \alpha, \gamma) + A_{z'}(y; \gamma, \beta).$$

For fixed z, z' , each of these three random variables equals its mode except with probability at most $2e$. Since $6e < 1$, all three are simultaneously equal to their modes for at least one choice of labels and y . Their identity gives $L(z + z') = L(z) + L(z')$. Thus L is a homomorphism.

Finally, (II.138.1) also says

$$e = \mathbb{P}_{u,v,\alpha,\beta,\gamma}[F(u, \alpha) \neq A_u(v; \gamma, \beta)].$$

Conditional on u, α , whenever $F(u, \alpha) \neq L(u)$, the event $A_u(v; \gamma, \beta) = L(u)$ forces a failure. This event has probability at least $1 - 2e$ and its labels are independent of α . Therefore $e \geq (1 - 2e)\mathbb{P}[F(u, \alpha) \neq L(u)]$, proving the claim. $\square$

Lemma II.138.2 (Noisy affine correction). *Let $F : G \times S \rightarrow K$. Let e be the probability that*

$$F(x_1, \alpha_1) + F(x_2, \alpha_2) \neq F(x_3, \alpha_3) + F(x_4, \alpha_4),$$

where $x_1 + x_2 = x_3 + x_4$ is uniformly sampled and all four labels are independent uniform elements of S . There is an affine map $A : G \rightarrow K$ such that

$$\mathbb{P}_{x,\alpha}[F(x, \alpha) \neq A(x)] \leq 6e. \quad (\text{II.138.4})$$

If $e < 1/6$, the stronger bound $e/(1 - 2e)$ holds.

Proof. For $e \geq 1/6$, any affine map gives (II.138.4). Suppose $e < 1/6$. Condition on $(x_4, \alpha_4) = (z, t)$ and choose an anchor whose conditional error e_0 is at most e . Conditional on this anchor, write $x_1 = z + u$, $x_2 = z + v$, $x_3 = z + u + v$, with u, v independent uniform. The function $Q(u, \alpha) = F(z + u, \alpha) - F(z, t)$ has noisy homomorphism error e_0 . Lemma II.138.1 gives a homomorphism L with error at most $e_0/(1 - 2e_0) \leq e/(1 - 2e)$. Translate back: $A(x) = L(x - z) + F(z, t)$ is affine. For $e < 1/6$, the stronger bound is at most $3e/2$, and in particular at most $6e$. $\square$

Lemma II.138.3 (Separation of affine maps). *If $A, A' : G \rightarrow K$ are distinct affine maps, then $\mathbb{P}_x[A(x) \neq A'(x)] \geq 1/2$.*

Proof. The difference has the form $L(x) + c$. If $L = 0$ and $c \neq 0$, it is nonzero everywhere. If $L \neq 0$, its zero set is either empty or a coset of $\ker L$. The index of that kernel is at least two. Thus its zero set has density at most $1/2$. $\square$

II.139 Synchronizing derivatives into one biadditive model

Source 33, delivered Section 4.

A separately chosen affine approximation to each derivative need not initially be compatible with addition of increments. The derivative cocycle and affine separation supply exactly the required compatibility.

Theorem II.139.1 (A coarse full-product decoder). *If $\varepsilon_4(\phi) < 1/216$, then*

$$\delta(\phi) \leq 6\varepsilon_4(\phi) + \frac{36\varepsilon_4(\phi)}{1 - 72\varepsilon_4(\phi)} \leq 60\varepsilon_4(\phi). \quad (\text{II.139.1})$$

Proof. Let e_h be the conditional failure probability of the $s = 4$ test at increment h . Apply Lemma II.138.2 to $F_h(x, y) = D_h\phi(x, y)$, with auxiliary-label set $S = H$. Choose $A_h \in \text{Aff}(G, K)$ such that

$$d_h := \mathbb{P}_{x,y}[D_h\phi(x, y) \neq A_h(x)] \leq 6e_h.$$

Set $d = \mathbb{E}_h d_h$. Then $d \leq 6\varepsilon_4(\phi) < 1/36$.

The exact derivative identity is

$$D_{h+k}\phi(x, y) = D_h\phi(x, y + k) + D_k\phi(x, y). \quad (\text{II.139.2})$$

If all three approximations hold at these three queries, then $A_{h+k}(x) = A_h(x) + A_k(x)$. Averaging the union bound gives

$$\mathbb{P}_{h,k,x}[A_{h+k}(x) \neq A_h(x) + A_k(x)] \leq 3d. \quad (\text{II.139.3})$$

For fixed h, k , either the two affine maps in this comparison are identical or their disagreement density is at least $1/2$, by Lemma II.138.3. Hence

$$\eta := \mathbb{P}_{h,k}[A_{h+k} \neq A_h + A_k] \leq 6d < 1/6.$$

The affine maps form a finite abelian group under pointwise addition. Lemma II.138.1, now with a singleton label set, gives a homomorphism

$$\mathcal{A} : H \longrightarrow \text{Aff}(G, K)$$

whose disagreement fraction with $h \mapsto A_h$ is at most

$$q \leq \frac{\eta}{1 - 2\eta} \leq \frac{6d}{1 - 12d}. \quad (\text{II.139.4})$$

Split each affine map into its constant and homomorphism parts by setting

$$\ell(y) = \mathcal{A}(y)(0), \quad B(x, y) = \mathcal{A}(y)(x) - \mathcal{A}(y)(0).$$

Because $\mathcal{A}$ is a homomorphism, ℓ is additive and B is additive in y ; because $\mathcal{A}(y)$ is affine, B is additive in x . For the residual $r(x, y) = \phi(x, y) - B(x, y) - \ell(y)$, we have

$$\mathbb{P}_{h,x,y}[D_h r(x, y) \neq 0] \leq d + q.$$

Indeed, outside the q -fraction of bad increments, disagreement is bounded by the original derivative error.

For fixed x , write $p_x(t) = \mathbb{P}_y[r(x, y) = t]$. The pair $(y, y + h)$ is independent uniform after averaging over y, h . Thus

$$\mathbb{P}_{h,y}[D_h r(x, y) \neq 0] = 1 - \sum_t p_x(t)^2 \geq 1 - \max_t p_x(t).$$

Choose $a(x)$ to be a most likely value of $r(x, y)$. The resulting model $a(x) + B(x, y) + \ell(y)$ has distance at most $d + q$. Using (II.139.4) and $d \leq 6\varepsilon_4$ proves the first inequality in (II.139.1). Finally, $1 - 72\varepsilon_4 > 2/3$, so the fraction in (II.139.1) is at most $54\varepsilon_4$. $\square$

Corollary II.139.2 (Exact classification). *For any even $s \geq 4$, $\varepsilon_s(\phi) = 0$ if and only if $\phi \in \mathcal{V}(G, H; K)$.*

Proof. Every model passes every test. Conversely, the moment monotonicity of Proposition II.137.2 gives $\varepsilon_4 \leq \varepsilon_s = 0$. Theorem II.139.1 then gives $\delta(\phi) = 0$. $\square$

Remark II.139.3 (Where the constants are spent). The coarse constant 60 is not the final repair coefficient. It merely puts a nearest model in a small neighborhood where a much more precise counting argument applies. The losses are explicitly localized: the factor 6 for derivative-wise affine correction, the three terms in the cocycle, affine separation by $1/2$, and one more majority correction. No density exponent or group-size dependence is hidden in this step.

II.140 Sharp local defect inequalities

Source 33, delivered Section 5.

II.140.1 The geometry of a nearest model

Fix a nearest model $\psi = a + B + \ell$ and let $r = \phi - \psi$. For fixed B, ℓ , each row constant can be optimized independently. We may therefore choose the nearest model so that zero is a most likely value of $r(x, y)$ in every row. Write

$$p_x(z) = \mathbb{P}_y[r(x, y) = z], \quad \delta_x = 1 - p_x(0), \quad \delta = \mathbb{E}_x \delta_x. \quad (\text{II.140.1})$$

Adding or subtracting an exact model does not change the outcome of any arrangement test. Thus $\varepsilon_s(r) = \varepsilon_s(\phi)$.

For each increment put

$$t_h = \mathbb{P}_{x,y}[D_h r(x, y) \neq 0], \quad \bar{t} = \mathbb{E}_h t_h. \quad (\text{II.140.2})$$

The following two simple inequalities capture the difference between pointwise corruption and concentrated corruption of a few rows.

Lemma II.140.1 (Derivative profile). *The residual of a nearest model satisfies*

$$0 \leq t_h \leq 2\delta \quad \text{for every } h, \quad \bar{t} = \mathbb{E}_x \left(1 - \sum_z p_x(z)^2 \right) \geq \delta. \quad (\text{II.140.3})$$

If in addition $\delta_x \leq \kappa$ for every x , then

$$\bar{t} \geq 2(1 - \kappa)\delta. \quad (\text{II.140.4})$$

Proof. A nonzero difference requires at least one of its two endpoint values to be nonzero, so $t_h \leq 2\delta$ by translation invariance and the union bound. Averaging over h, y gives two independent uniform points of each row and therefore the collision formula in (II.140.3). Since $p_x(0) = \max_z p_x(z)$, $\sum_z p_x(z)^2 \leq p_x(0)$, proving $\bar{t} \geq \delta$. Also

$$\sum_z p_x(z)^2 \leq (1 - \delta_x)^2 + \delta_x^2,$$

because the total mass outside zero is δ_x . Hence the row's noncollision probability is at least $2\delta_x(1 - \delta_x)$, which is at least $2(1 - \kappa)\delta_x$ under the extra hypothesis. $\square$

II.140.2 Exactly one erroneous edge

Lemma II.140.2 (One-error probability with one missing independence degree). *Let $s \geq 4$ be even. Suppose $Z_1, \dots, Z_s$ are Bernoulli random variables with common mean t , and every proper subfamily is independent. Then*

$$\mathbb{P}\left[\sum_{i=1}^s Z_i = 1\right] \geq L_s(t) := st(1-t)((1-t)^{s-2} - t^{s-2}). \quad (\text{II.140.5})$$

The ratio $L_s(t)/t$, extended continuously at zero, is nonnegative and nonincreasing on $[0, 1/2]$.

Proof. Let $p_* = \mathbb{P}[Z_1 = \dots = Z_s = 1]$. Independence of the first $s-1$ variables gives $p_* \leq t^{s-1}$. Inclusion-exclusion at a specified position i gives

$$\mathbb{P}[Z_i = 1, Z_j = 0 \ (j \neq i)] = t(1-t)^{s-1} + t^s - p_*.$$

Here the last sign is negative because s is even. Summing over i and using the upper bound on p_* yields

$$\begin{aligned} \mathbb{P}\left[\sum_i Z_i = 1\right] &\geq st(1-t)^{s-1} + st^s - st^{s-1} \\ &= st(1-t)((1-t)^{s-2} - t^{s-2}). \end{aligned}$$

For $0 \leq t \leq 1/2$ this expression is nonnegative. After division by t , its derivative, divided by s , equals

$$-(s-1)(1-t)^{s-2} - (s-2)t^{s-3} + (s-1)t^{s-2},$$

which is nonpositive on that interval. The extension at zero is s . $\square$

Theorem II.140.3 (Profile-sensitive local bound). *Let $0 < \delta(\phi) \leq 1/4$, and use the collision profile of any nearest model chosen as in (II.140.1). Then*

$$\varepsilon_s(\phi) \geq \frac{\bar{t}}{\delta} F_s(\delta) \geq F_s(\delta). \quad (\text{II.140.6})$$

If every row error is at most $\kappa \leq 1/2$, this strengthens to

$$\varepsilon_s(\phi) \geq 2(1 - \kappa)F_s(\delta). \quad (\text{II.140.7})$$

Proof. At fixed h , let $Z_i = \mathbf{1}_{\{D_h r(x_i, y_i) \neq 0\}}$. By Lemma II.137.1, these variables satisfy the hypotheses of Lemma II.140.2, with mean t_h . If exactly one difference is nonzero, the signed sum of differences is nonzero in K , so the arrangement fails. Therefore its conditional failure probability is at least $L_s(t_h)$.

Because $0 \leq t_h \leq 2\delta \leq 1/2$, monotonicity of $L_s(t)/t$ gives

$$\varepsilon_s(\phi) \geq \mathbb{E}_h L_s(t_h) \geq \frac{L_s(2\delta)}{2\delta} \mathbb{E}_h t_h = \frac{\bar{t}}{\delta} F_s(\delta).$$

Apply Lemma II.140.1. The case $\delta = 0$, omitted from the ratio notation, has $\varepsilon_s = 0$ and satisfies the unprofiled bound as well. $\square$

This proves Theorem F.38.3, except for its equality statement, which follows from the explicit family in Section II.141. It also explains the sharp leading constant. If errors are uniformly sparse in every row, with $\kappa = o(1)$, the leading lower bound approaches $2s\delta$, not $s\delta$. The worst examples for the general theorem instead concentrate errors in a small set of rows with substantial internal disagreement.

II.140.3 A quadratic estimate and the global small branch

The full expression F_s is useful for local information, but a simpler quadratic has a transparent inverse.

Lemma II.140.4 (Quadratic local defect). *If $\delta(\phi) \leq 1/(2(s-1))$, then*

$$\varepsilon_s(\phi) \geq Q_s(\delta) := s\delta - 2s(s-1)\delta^2. \quad (\text{II.140.8})$$

Proof. At fixed h , let Z_i be as above and set $Z = \sum_i Z_i$. For every nonnegative integer Z , $\mathbf{1}_{\{Z=1\}} \geq Z - Z(Z-1)$. Pairwise independence implies

$$\mathbb{P}[Z=1] \geq st_h - s(s-1)t_h^2.$$

Using $t_h \leq 2\delta$ and then $\bar{t} \geq \delta$ yields

$$\varepsilon_s \geq s(1 - 2(s-1)\delta)\bar{t} \geq s(1 - 2(s-1)\delta)\delta.$$

The last inequality uses the stated upper bound on δ to ensure that the multiplier is nonnegative. $\square$

Proof of the repair statement in Theorem F.38.2. Put $\varepsilon = \varepsilon_s(\phi)$. Proposition II.137.2 and the coarse decoder give

$$\delta \leq 60\varepsilon_4 \leq 60\varepsilon \leq \frac{1}{4(s-1)}.$$

The coarse decoder applies because $1/(240(s-1)) \leq 1/720 < 1/216$. The function Q_s is increasing on $[0, 1/(4(s-1))]$. Lemma II.140.4 gives $Q_s(\delta) \leq \varepsilon$, and the smaller root of $Q_s(v) = \varepsilon$ is exactly $R_s(\varepsilon)$. Its discriminant is positive in the stated range. Thus $\delta \leq R_s(\varepsilon)$. Taylor expansion at zero proves (F.38.6). $\square$

Remark II.140.5 (Why the coarse step cannot be skipped). A local defect inequality alone need not imply a small distance from small failure probability: its right side may fall again at larger δ . In particular, $F_4(1/4) = 0$. The coarse decoder chooses the correct inverse branch before we use a sharper polynomial. Omitting this step would leave a genuine logical gap.

II.140.4 The cubic inverse for four differences

Corollary II.140.6 (Cubic repair modulus). *For $0 \leq u \leq 1/720$, let $\rho_4(u)$ be the unique root in $[0, 1/12]$ of*

$$4v - 24v^2 + 32v^3 = u. \quad (\text{II.140.9})$$

If $\varepsilon_4(\phi) \leq 1/720$, then

$$\delta(\phi) \leq \rho_4(\varepsilon_4(\phi)), \quad (\text{II.140.10})$$

$$\rho_4(u) = \frac{u}{4} + \frac{3u^2}{8} + u^3 + \frac{105u^4}{32} + O(u^5). \quad (\text{II.140.11})$$

The entire cubic inequality, and hence this inverse at an infinite sequence of errors tending to zero, is attained by explicit examples.

Proof. The cubic is strictly increasing on $[0, 1/12]$: its derivative $4 - 48v + 96v^2$ is positive there. Its endpoint values are 0 and $5/27$, so the specified root exists and is unique. The coarse decoder gives $\delta \leq 60\varepsilon_4 \leq 1/12$, and the local cubic bound gives (II.140.10). Substitution of a power series in (II.140.9) gives (II.140.11); equivalently, the analytic inverse theorem applies since the derivative at zero is four. Equality follows from Theorem II.141.1 below for all sufficiently large odd n . $\square$

For comparison, the quadratic inverse at $s = 4$ starts with $u/4 + 3u^2/8 + 9u^3/8$. The sharp cubic reduces its third coefficient from $9/8$ to 1. Sharpness here does not mean that every real number u is a feasible finite-group failure probability; it means that the inequality is attained at the explicitly specified finite examples, including examples arbitrarily close to zero.

II.141 Extremal families and optimal coefficients

Source 33, delivered Section 6.

II.141.1 An exact binary family

Theorem II.141.1 (Sharp row-defect family). *Let $n \geq 3$ be odd, and define*

$$\phi_n : \mathbb{Z}_n \times \mathbb{Z}_2 \longrightarrow \mathbb{Z}_2, \quad \phi_n(x, y) = \mathbf{1}_{\{0\}}(x)y.$$

Then

$$\delta_n := \delta(\phi_n) = \frac{1}{2n}, \quad (\text{II.141.1})$$

and for every even $s \geq 4$,

$$\varepsilon_s(\phi_n) = \frac{1}{4} \left[1 - \left(1 - \frac{2}{n} \right)^s - (n-1) \left(\frac{2}{n} \right)^s \right]. \quad (\text{II.141.2})$$

For $s = 4$ this is exactly

$$\varepsilon_4(\phi_n) = \frac{2}{n} - \frac{6}{n^2} + \frac{4}{n^3} = 4\delta_n - 24\delta_n^2 + 32\delta_n^3. \quad (\text{II.141.3})$$

Proof. Every homomorphism $\mathbb{Z}_n \rightarrow \mathbb{Z}_2$ is zero, since n is odd. Consequently every biadditive map $\mathbb{Z}_n \times \mathbb{Z}_2 \rightarrow \mathbb{Z}_2$ is zero, and the exact models are $a(x) + cy$ with $c \in \mathbb{Z}_2$. The row at $x = 0$ has slope one; the other $n - 1$ rows have slope zero. If a row has the wrong slope, changing its constant can match exactly one of its two values. The best choice is $c = 0$, with just one wrong-slope row, and its distance is $1/(2n)$.

The increment $h = 0$ always passes. At $h = 1$, the derivative is $\mathbf{1}_{\{0\}}(x)$, independent of y . The nontrivial character of $\mathbb{Z}_2$ produces $f(x) = (-1)^{\mathbf{1}_{\{0\}}(x)}$. Its normalized Fourier coefficients on $\mathbb{Z}_n$ are $1 - 2/n$ at the trivial character and $-2/n$ at every nontrivial character. By character orthogonality, conditional failure at $h = 1$ is

$$\frac{1}{2} \left[1 - \sum_{\xi \in \widehat{\mathbb{Z}_n}} |\widehat{f}(\xi)|^s \right].$$

Average over the two increments to obtain (II.141.2). Expanding at $s = 4$ proves (II.141.3). $\square$

The equality is not a rounding artifact. For this family the residual has error $1/2$ in exactly one row, so $\bar{t} = \delta_n$; $t_0 = 0$ and $t_1 = 2\delta_n$. At $s = 4$, the all-four-error event has probability $n^{-3} = t_1^3$, saturating the upper bound in Lemma II.140.2. Three zero horizontal coordinates force the fourth to be zero. Thus a three-error event is impossible, and the binary test fails precisely on a one-error event. All the inequalities in the local cubic proof become equalities.

Corollary II.141.2 (Sharp first two asymptotic coefficients). *Uniformly over finite abelian groups, the coefficients $1/s$ and $2(s - 1)/s^2$ in (F.38.6) are optimal in succession.*

Proof. For fixed even $s \geq 4$, expansion of (II.141.2), with $\delta_n = 1/(2n)$, gives

$$\varepsilon_s(\phi_n) = s\delta_n - 2s(s - 1)\delta_n^2 + O_s(\delta_n^3).$$

The term $(n - 1)(2/n)^s$ first contributes at order δ_n^{s-1} , which is at least three. Inverting the displayed expansion gives

$$\delta_n = \frac{\varepsilon_s(\phi_n)}{s} + \frac{2(s - 1)}{s^2} \varepsilon_s(\phi_n)^2 + O_s(\varepsilon_s(\phi_n)^3).$$

These are actual examples with errors tending to zero. They rule out a smaller universal leading coefficient, and, with that leading coefficient fixed, a smaller universal quadratic coefficient. For $s = 4$, (II.141.3) also attains every coefficient of the analytic small-branch cubic inverse along this sequence. $\square$

II.141.2 The leading constant in Gowers's common prime setting

The preceding family's second-order sharpness uses two-torsion, but its leading-order obstruction is not an artifact of characteristic two. We now keep all three groups equal to the same odd prime field.

Theorem II.141.3 (Prime cyclic sharpness of the leading constant). *For an odd prime $p \geq 5$, define*

$$\phi_p : \mathbb{F}_p \times \mathbb{F}_p \rightarrow \mathbb{F}_p, \quad \phi_p(x, y) = \mathbf{1}_{\{0\}}(x)y.$$

Then

$$\delta_p = \frac{p - 1}{p^2}. \tag{II.141.4}$$

Fix even $s = 2m \geq 4$. As $p \rightarrow \infty$ through primes,

$$\varepsilon_s(\phi_p) = s\delta_p - \frac{s(3s-2)}{4}\delta_p^2 + O_s(\delta_p^3). \quad (\text{II.141.5})$$

In particular, the coefficient $1/s$ in the repair bound is optimal already for common odd prime cyclic groups.

Proof. A model has the form $a(x) + (bx+c)y$. If its slope at a row agrees with $\mathbf{1}_{\{0\}}(x)$, an appropriate row constant matches all p values; otherwise it matches exactly one. The zero affine slope agrees on $p-1$ rows. A nonconstant affine slope can agree with the zero slope on at most one nonzero row and with slope one at $x=0$ on at most one additional row, hence on at most two rows. Constant nonzero slopes do no better. Since $p \geq 5$, the zero affine slope is optimal and gives (II.141.4).

For the arrangement error, put $\nu = 1/p$ and assume $p > m$, which is sufficient for the asymptotic claim. The derivative is $D_h\phi_p(x, y) = h\mathbf{1}_{\{0\}}(x)$. Let $Z_i = \mathbf{1}_{\{0\}}(x_i)$. Every proper subfamily of the Z_i is independent Bernoulli(ν), whereas

$$\mathbb{P}[Z_1 = \cdots = Z_s = 1] = \nu^{s-1}.$$

Inclusion–exclusion shows that for any subset I of positions, the probability that its positions are exactly the ones with $Z_i = 1$ is

$$\nu^{|I|}(1-\nu)^{s-|I|} + (-1)^{s-|I|}(\nu^{s-1} - \nu^s). \quad (\text{II.141.6})$$

For $h \neq 0$, the test accepts if the number of ones in the first half equals the number in the second half modulo p . Since $p > m$, this is ordinary equality of the two integer counts. Summing (II.141.6) over the accepting patterns and then averaging over h gives the exact expression

$$\begin{aligned} \varepsilon_s(\phi_p) = (1-\nu) \left[1 - \sum_{a=0}^m \binom{m}{a}^2 \nu^{2a} (1-\nu)^{s-2a} \right. \\ \left. - \binom{s}{m} (\nu^{s-1} - \nu^s) \right]. \end{aligned} \quad (\text{II.141.7})$$

We used $\sum_a \binom{m}{a}^2 = \binom{2m}{m}$. The sum in brackets expands as

$$s\nu - \left(\binom{s}{2} + m^2 \right) \nu^2 + O_s(\nu^3) = s\nu - \frac{s(3s-2)}{4} \nu^2 + O_s(\nu^3).$$

The last correction in (II.141.7) starts at order $s-1 \geq 3$. Since $\delta_p = (1-\nu)\nu$, this proves (II.141.5) and the leading-coefficient claim. $\square$

The prime family also shows that, with leading coefficient $1/s$ fixed, a repair estimate valid for all sufficiently large common prime fields cannot have quadratic coefficient smaller than

$$\frac{3s-2}{4s^2}. \quad (\text{II.141.8})$$

Our general upper bound has coefficient $2(s-1)/s^2$. Determining the optimal coefficient between these two values in the common-prime setting is a concrete unresolved refinement of this work, not an asserted consequence of the binary example.

II.141.3 What is, and is not, sharp

The following distinctions prevent several possible overinterpretations. The whole polynomial F_4 is attained at the distances in (II.141.1). The global inverse statement uses only its increasing small branch. For $s \geq 6$, the full polynomial F_s is not asserted sharp; the binary example and the proved lower bound need not have the same cubic term. The all-group quadratic inverse coefficient is sharp, while only the leading coefficient has been proved sharp here under the common odd-prime restriction. Finally, none of the convenient numerical activation thresholds is claimed optimal.

II.142 Erasures and almost-full products

Source 33, delivered Section 7.

II.142.1 Repair on a partial domain

Let $A \subset G \times H$ have density $1 - \tau > 0$, and let $\phi : A \rightarrow K$. Draw a full-domain arrangement as before and let V_A be the event that all $2s$ queried point occurrences lie in A . Write $\lambda = \mathbb{P}[V_A]$. Each point occurrence is uniform on $G \times H$, so

$$\lambda \geq \max\{0, 1 - 2s\tau\}. \quad (\text{II.142.1})$$

No independence of the $2s$ point occurrences is needed for this union bound.

Theorem II.142.1 (Erasure-robust repair). *Suppose $\lambda > 0$ and, conditional on V_A , the arrangement failure probability is at most $\eta \in [0, 1]$. Set*

$$E = \lambda\eta + (1 - \lambda). \quad (\text{II.142.2})$$

If $E \leq 1/(240(s - 1))$, then there is a model $\psi \in \mathcal{V}(G, H; K)$ with

$$\frac{|\{(x, y) \in A : \phi(x, y) \neq \psi(x, y)\}|}{|A|} \leq \frac{R_s(E)}{1 - \tau}. \quad (\text{II.142.3})$$

A sufficient condition, not requiring the exact value of λ , is $\eta + 2s\tau \leq 1/(240(s - 1))$; in that case the right side may use $R_s(\eta + 2s\tau)/(1 - \tau)$.

Proof. Extend ϕ arbitrarily to a function $\tilde{\phi} : G \times H \rightarrow K$. Its test agrees with the partial-domain test on V_A , and every other arrangement may pessimistically be counted as a failure. Therefore $\varepsilon_s(\tilde{\phi}) \leq E$. Apply Theorem F.38.2; monotonicity of R_s on the relevant interval gives an absolute disagreement density at most $R_s(E)$. Restricting the set of disagreements to A and normalizing by $|A| = (1 - \tau)|G||H|$ proves (II.142.3). Finally,

$$E = \eta + (1 - \eta)(1 - \lambda) \leq \eta + 2s\tau$$

by (II.142.1). Under the sufficient condition, $1 - 2s\tau > 0$, so validity is automatically nonvacuous. $\square$

For eight-arrangements the simple effective error is $\eta + 32\tau$. This is a near-full-domain result. If A has small density, the upper bound for invalid arrangements is too large to activate the decoder; claiming a useful estimate in that regime would be unjustified.

II.142.2 Exact counting in the presence of erasures

For cyclic groups define

$$R_h(x) = |\{y : (x, y), (x, y + h) \in A\}|,$$

$$M_h^A(x, z) = |\{y : (x, y), (x, y + h) \in A, D_h\phi(x, y) = z\}|.$$

Using convolution on $\mathbb{Z}_n$ in the first expression and on $\mathbb{Z}_n \times \mathbb{Z}_q$ in the second, respectively, gives

$$N_{\text{valid}} = \sum_{h,x} (R_h^{*(s/2)}(x))^2, \quad (\text{II.142.4})$$

$$N_{\text{respected}} = \sum_{h,x,z} ((M_h^A)^{*(s/2)}(x, z))^2. \quad (\text{II.142.5})$$

The proof is the same matching-of-two-halves argument as for (II.137.8). Thus both λ and the conditional failure probability are exact rational numbers. If $N_{\text{valid}} = 0$, the conditional probability is undefined rather than silently interpreted as zero.

II.143 Exact cyclic decoding and reproducible checks

Source 33, delivered Section 8.

II.143.1 A nearest-model algorithm

The model class contains an arbitrary function a , but it need not be searched exponentially. For fixed bilinear and linear coefficients, its rows separate.

Proposition II.143.1 (Exact cyclic decoder). *Let $\phi : \mathbb{Z}_n \times \mathbb{Z}_m \rightarrow \mathbb{Z}_q$. Every exact model has the form*

$$\psi(x, y) = a_x + (bx + c)y \pmod{q}, \quad nb \equiv mb \equiv mc \equiv 0 \pmod{q}. \quad (\text{II.143.1})$$

There are $g_b = \gcd(n, m, q)$ choices of b and $g_c = \gcd(m, q)$ choices of c . For each such pair, choose a_x as a mode of the offsets

$$\phi(x, y) - (bx + c)y \pmod{q} \quad (y \in \mathbb{Z}_m). \quad (\text{II.143.2})$$

The best pair and these modal offsets give an exact nearest model. The search uses $O(nmg_bg_c)$ modular and histogram operations.

Proof. A biadditive map is determined by $b = B(1, 1)$, so it must be bxy ; well-definedness is exactly $nb = mb = 0$ in $\mathbb{Z}_q$. Similarly a homomorphism in y has the form cy with $mc = 0$. The counts of solutions to these annihilation constraints are the stated greatest common divisors. For fixed b, c , maximizing matches in row x is precisely choosing a most frequent value in (II.143.2). Rows are independent in this optimization, so their maximal scores add. Enumerate the g_bg_c pairs and retain the largest total score. $\square$

The implementation uses sparse Python histograms with deterministic tie breaking; the operation count treats dictionary access in its usual expected-constant-time model. It does not claim the same

number of bit operations on unbounded integers. In the common prime case the candidate search is $O(p^4)$ arithmetic/histogram operations.

For exact arrangement counting, binary exponentiation of convolution in (II.137.8) uses $O(m(nq)^2 \log s)$ integer arithmetic operations with a direct dense convolution bound. The supplied implementation skips zero entries. This is a verification algorithm, not a claim of an asymptotically best large-scale counter. Coefficients are nonnegative Python integers, and probabilities use `fractions.Fraction`.

II.143.2 The shipped verification suite

The file `code/verify.py` checks the identities and inequalities against exact finite examples. The default suite exhausts the first five families of Table II.143.1; the `-extended` option adds the sixth. The recorded delivery run used that option.

n	m	q	Tables exhausted	Zero-error tables	Literal check
2	2	2	16	16	yes
3	2	2	64	16	yes
3	3	2	512	8	no
2	3	3	729	27	yes
3	4	2	4,096	16	no
3	3	3	19,683	243	no
Total			25,100		

Table II.143.1: Exact full-table enumeration. Every table was tested at $s = 4, 6, 16$. The last column means that the $s = 4$ counter and nearest distance were also compared with independent literal enumerations.

For each full table the suite checks exact classification, both cross-order defect bounds, the local bound F_s , the stronger collision-profile bound, the quadratic inequality when applicable, and the coarse and small-branch estimates when their hypotheses hold. The smallest groups were additionally checked against direct enumeration of every arrangement and of every full model, including all row constants. This independently checks both the convolution counter and the modal decoder.

There are also 300 seeded sample tables, 100 each for $(n, m, q) = (4, 4, 4), (5, 4, 6), (5, 5, 5)$, using seed 20261006. The binary family is compared with the exact counter for $n = 3, 5, 7, 9, 11$, and its formula is checked at larger values, including $n = 100001$, where the high-agreement estimates are nonvacuous. The common-prime row family is compared with exact counts for $p = 3, 5, 7, 11, 13, 17$. The implementation of its formula includes congruence wraparound and so also handles primes not covered by the simplified formula (II.141.7).

A separate erasure suite exhausts all $3^6 = 729$ partial tables on $\mathbb{Z}_3 \times \mathbb{Z}_2 \rightarrow \mathbb{Z}_2$, allowing each entry to be zero, one, or missing. It checks (II.142.4)–(II.142.5) by literal enumeration at $s = 4$, tests $s = 4, 6, 16$ and the ambient-normalized moment bounds, and checks both constant extensions against the effective-error and validity bounds. Only the fully erased table has no valid arrangement; no conditional failure probability is assigned to it.

Negative tests reject malformed tables, invalid orders and moduli, and an intentionally corrupted model score on every full-table test. A separate small rational-polynomial module verifies the cubic reversion coefficients and the first two sharp-family coefficients. There is no floating-point

optimization or acceptance of a numerical solver’s feasibility status in these checks.

Example II.143.2 (An exact reproducible input). The included `data/example.json` is the binary family at $n = 5$. Its nearest-model distance is $1/10$. The four-difference test has 3232 respected arrangements out of 4000, hence failure $24/125$. The eight-arrangement failure probability is $7627229088/30517578125$. These values are not in the small-error activation range; they check the local formula and normalization, not a purported global inverse outside its hypotheses.

II.143.3 What computation establishes

The checks verify the listed finite claims and expose errors in normalization, sign conventions, implementation, and polynomial arithmetic. They do not prove the theorems for arbitrary groups, certify literature novelty, or replace the arguments in Sections II.138–II.141. Most exhaustive small-group tables have either zero error or an error larger than the conservative activation threshold. The explicit large sharp-family calculations, rather than those small enumerations, are what exercise nonzero inputs inside that threshold. The general proof of the threshold remains analytic and combinatorial.

II.144 Source 35’s Part 1: degree-weighted selection at Lemma 15.5 [write]

Added in the seventh write, 6 October 2026.

The next four sections are source 35’s Sections 2–5, its delivered Part 1, on the arrangement-selection step of Gowers’s Lemma 15.5: retain many d -arrangements of k -cubes while making almost all retained arrangements respected. *Logically the chapter follows the Section 15 material* of sources 17 (its Fejér kernel), 25 (selection for every $d \geq 2$) and 27 (degeneracy counts), and source 40’s one-stage Fejér selection, which is placed in this directory and not yet written in; it is printed here because no number printed earlier may move. Source 35’s introduction is Section F.40.

[write] **7 October 2026; sources 40–46.** *The forward reference is resolved:* source 40 is now printed, at the end of Part II (Section II.164); logically it comes before this chapter. Its one-stage Fejér selection has field-size exponent $2s$ and retention $\alpha^s \eta^{s-1} / (3^{s-1} (s+1)^s)$, which this chapter improves (note after source 40’s Theorem II.169.2).

What is proved. Coefficient arrays are sampled with a product law whose largest atom is $1/L$, and an array of exact degree t (after removing Boolean parity) must lie in a balance space of dimension $D_t = m \sum_{i \leq t} \binom{k}{i} - \binom{k}{t}$ while imposing at least $t + 1$ independent restrictions; so the sufficient field-size exponent is $E_{k,m} = \max_t D_t / (t + 1) \sim 2s/k$ (23 against source 40’s $2s = 128$ at $d = 8, k = 2$), with retention $\alpha^s \eta^{s-1} / (2(s+2)^s)$ and purity $1/(1 + \eta)$, physical coincidences of vertices charged separately (Theorem II.145.2, Corollary II.145.3); a substitution at the Lemma 15.6 interface with the corrected exponents of Lemma 14.8 (Corollary II.148.3). The degree-sensitive certificate geometry is source 22’s, reproved for arbitrary field coefficients; the selector and the retained powers are source 40’s.

Relation to the formal statements (checked at HEAD 89c1cbc38). `lemma_15_5` (Sections14_15.lean, line 414) has the exact companion `lemma_15_5_holds` (Proofs15Restriction.lean, line 3116), with an existential threshold; whether Corollary II.145.3 gives it with an explicit threshold is a marked question (item II.50). Nothing in this chapter is formalized.

II.145 The arrangement model and the quantitative theorem

Source 35, delivered Section 2.

Delivered part heading of source 35: **Degree-weighted arrangement selection**

Throughout this part let $k \geq 1$, $d \geq 2$, $m = 2d$, and

$$r = 2^k, \quad s = mr, \quad M = (m+1)k + m - 1.$$

Let $\mathbb{F}_q$ have characteristic p , and put

$$\sigma = (\underbrace{1, \dots, 1}_d, \underbrace{-1, \dots, -1}_d), \quad W = \{z \in \mathbb{F}_q^m : \sigma \cdot z = 0\}.$$

For $e \in \{0, 1\}^k$, write $\pi(e) = (-1)^{|e|}$. The product $e \odot h$ means coordinatewise multiplication.

Definition II.145.1 (Ordered arrangements and their respected relation). An arrangement parameter is a tuple

$$(h, x_1, \dots, x_m, z) \in \mathbb{F}_q^k \times (\mathbb{F}_q^k)^m \times W.$$

Its labelled vertices are

$$v_{j,e} = (x_j + e \odot h, z_j) \in \mathbb{F}_q^{k+1}, \quad 1 \leq j \leq m, \quad e \in \{0, 1\}^k.$$

The parameters, including degenerate ones, are counted with their given order. There are exactly q^M parameters. An arrangement lies in $B \subseteq \mathbb{F}_q^{k+1}$ if all its vertices belong to B . A map $\phi : B \rightarrow \mathbb{F}_q$ respects it if

$$\sum_{j=1}^m \sum_{e \in \{0,1\}^k} \sigma_j \pi(e) \phi(v_{j,e}) = 0. \quad (\text{II.145.1})$$

Write $T(B)$ for the number of arrangements in B , and $G_\phi(B)$ for the number it respects.

If $|B| = \beta q^{k+1}$, then

$$T(B) \leq T_0, \quad T_0 = \beta^{m-1} q^M. \quad (\text{II.145.2})$$

Indeed, choose h , choose the first $m-1$ origin vertices in B , and choose the last horizontal anchor in $\mathbb{F}_q^k$. The relation on z determines its last level. This gives the upper bound $q^{2k}|B|^{m-1}$, which is exactly T_0 . This normalization agrees with the arrangement interface in Gowers' selection argument [1, 158].

For $0 \leq t \leq k$ define

$$c_t = \sum_{i=0}^t \binom{k}{i}, \quad D_t = mc_t - \binom{k}{t}. \quad (\text{II.145.3})$$

In particular, $D_0 = m-1$ and $D_k = s-1$. The probability factors needed below are

$$Q_0(q) = q^{-1}, \quad Q_1(q) = q^{-2} + (q-1)q^{-m}, \quad Q_t(q) = q^{-t-1} \quad (2 \leq t \leq k). \quad (\text{II.145.4})$$

Finally put $C_{\text{coll}} = \binom{m}{2} 3^k$. The subscript refers to coincidences between physical vertices, not to linear dependencies between different coefficient arrays.

Theorem II.145.2 (Degree-weighted Fejér selection). *Suppose $0 < \alpha, \beta, \eta \leq 1$, $|B| = \beta q^{k+1}$, and*

$$G_\phi(B) \geq \alpha T_0.$$

Choose an integer $L \geq 2$, and assume $p > 2(L-1)$. Define

$$R_s(L) = 1 + 2 \sum_{a=1}^{L-1} (1 - a/L)^s, \quad (\text{II.145.5})$$

$$\mathcal{K}(L, q) = \sum_{t=0}^k L^{D_t} Q_t(q) + C_{\text{coll}} L^{s-1} q^{-k-1}, \quad (\text{II.145.6})$$

$$F_* = \alpha \eta R_s(L) - (1 - \alpha) - \beta^{-(m-1)} \mathcal{K}(L, q). \quad (\text{II.145.7})$$

If $F_ > 0$, there is $B' \subseteq B$ such that*

$$G_\phi(B') \geq \frac{F_*}{\eta L^s} T_0, \quad (\text{II.145.8})$$

$$\frac{G_\phi(B')}{T(B')} \geq \frac{1}{1 + \eta} \geq 1 - \eta, \quad (\text{II.145.9})$$

$$\frac{|B'|}{q^{k+1}} \geq \beta \left(\frac{F_*}{\eta L^s} \right)^{1/(m-1)}. \quad (\text{II.145.10})$$

[write] **6 October 2026; sources 34–39.** *Credits and comparisons* (checked at intake and at the write). The one-stage Fejér selector and the retained powers $\alpha^s \eta^{s-1}$ are source 40's (placed, not yet written; its field-size exponent is $2s$); the degree-sensitive certificate geometry is source 22's, reproved here for arbitrary field coefficients, as the source says (Section V.132). Source 17's Fejér kernel (its Section 9) and source 25's selection theorem for every $d \geq 2$ are neighbours, not duplicates. The table of $E_{k,m}$ was recomputed at intake: $E_{2,16} = \max\{15, 23, 21\} = 23$ against $2s = 128$; $E_{1,m} = (s-1)/2$ exactly, so the degree stratification gains only for $k \geq 2$; $E_{k,m}/(2s/k) = 0.57, 0.66, 0.73, 0.79$ at $m = 4$ and $k = 10, 20, 40, 80$ ($0.59, 0.67, 0.74, 0.79$ at $m = 16$; recomputed at the write), so the asymptotic $2s/k$ is approached slowly.

The formula retains the exact Fejér signal $R_s(L)$, the degree-one exceptional probability, and the separate collision cost. It can be used to optimize L for particular data. The following simpler consequence sacrifices that flexibility to give one explicit sufficient field-size condition.

Corollary II.145.3 (A convenient explicit threshold). *Under the density and respect hypotheses of Theorem II.145.2, set*

$$L = \left\lceil \frac{s+1}{\alpha \eta} \right\rceil, \quad a_t = 1 + \mathbf{1}_{t=1} + C_{\text{coll}} \mathbf{1}_{t=k}. \quad (\text{II.145.11})$$

Assume $p > 2(L-1)$ and

$$q \geq \max_{0 \leq t \leq k} \left(2(k+1) a_t \beta^{-(m-1)} L^{D_t} \right)^{1/(t+1)}. \quad (\text{II.145.12})$$

Then a subset $B' \subseteq B$ has purity at least $1/(1+\eta)$ and satisfies

$$G_\phi(B') \geq \frac{T_0}{2\eta L^s} \geq \frac{\alpha^s \eta^{s-1}}{2(s+2)^s} T_0. \quad (\text{II.145.13})$$

The density bound in (II.145.10) holds with F_ replaced by $1/2$. If $q = p$ is prime, condition (II.145.12) already implies the required characteristic condition.*

Deduction from the exact theorem. Since $m \geq 4$, $Q_1(q) \leq 2q^{-2}$. The collision term has degree weight $D_k = s - 1$. Consequently

$$\mathcal{K}(L, q) \leq \sum_{t=0}^k a_t L^{D_t} q^{-t-1} \leq \frac{\beta^{m-1}}{2}$$

by (II.145.12). Convexity of $u \mapsto (1 - u)^s$ on $[0, 1]$ shows that its trapezoidal sum is at least its integral. Therefore

$$R_s(L) \geq \frac{2L}{s+1}, \quad \alpha\eta R_s(L) \geq 2. \quad (\text{II.145.14})$$

Thus $F_* \geq 2 - (1 - \alpha) - 1/2 = 1/2 + \alpha > 1/2$. The exact theorem gives the first inequality of (II.145.13). Because $L \leq (s+2)/(\alpha\eta)$, it gives the second. For a prime field, the $t = 0$ condition implies $p \geq 2(k+1)\beta^{-(m-1)}L^{m-1} \geq 4L^{m-1} > 2(L-1)$. $\square$

[write] **6 October 2026; sources 34–39.** *Against the formal statement.* lemma_15_5 (Sections14_15.lean, line 414, at HEAD 89c1cbc38) has the exact companion lemma_15_5_holds (Proofs15Restriction.lean, line 3116), with an existential threshold N_0 and retention $(\alpha\eta/4)^{2^{2^{k+4}+k+3}}\beta^{15}N^{17k+15}$. Its count normalization $\beta^{15}N^{17k+15}$ is this corollary's $\beta^{m-1}q^M$ at $m = 16$, $q = N$ prime. The intake observed that this corollary then gives the formal conclusion with an explicit N_0 and retention $\alpha^s\eta^{s-1}/(2(s+2)^s)$, provided the ordered and degenerate conventions of `respectedGeneralArrangementCount` match those used here; that matching was not checked at the intake or at the write, so the observation is recorded as a marked question (item II.50). The formal statement is already exact, so nothing formal depends on it.

II.146 Degree balances and the probability of a certificate

Source 35, delivered Section 3.

II.146.1 Product probability on a subspace

The following elementary fact replaces a count of all available Fourier coefficients by the mass that the selector assigns to them.

Lemma II.146.1 (Subspace mass under a product law). *Let $X_1, \dots, X_n$ be independent random variables in $\mathbb{F}_q$, each with largest atom at most μ . If $V \subseteq \mathbb{F}_q^n$ is an affine subspace of dimension D , then*

$$\Pr((X_1, \dots, X_n) \in V) \leq \mu^{n-D}. \quad (\text{II.146.1})$$

More generally, if the i th largest atom is at most μ_i and projection onto a coordinate set I , $|I| = D$, is injective on V , the right side can be replaced by $\prod_{i \notin I} \mu_i$.

Proof. Choose D coordinates for which the coordinate projection on the direction space has rank D . Every choice of their values has at most one completion in V . Conditional on those values, independence bounds the probability of that completion by the product of the largest atoms of the other coordinates. Averaging over the chosen coordinates proves both claims. $\square$

The exponent in this general lemma cannot be increased for arbitrary subspaces: a coordinate subspace fixing $n - D$ coordinates at their largest atoms attains it. This does not assert sharpness for the particular arrangement subspaces below.

II.146.2 The degree of a coefficient array

Consider a field coefficient array $b = (b_{j,e}) \in \mathbb{F}_q^s$. For each cube, interpolate the values $\pi(e)b_{j,e}$ by the unique multilinear polynomial

$$V_j(u) = \sum_{T \subseteq [k]} A_{j,T} u_T, \quad u_T = \prod_{i \in T} u_i. \quad (\text{II.146.2})$$

Here and below the empty product is one. For $b \neq 0$, its parity-removed degree is

$$t(b) = \max\{|T| : A_{j,T} \neq 0 \text{ for some } j\}.$$

This definition is valid over every field: evaluation of multilinear polynomials on the Boolean cube is invertible by triangular interpolation.

For $S \subseteq [k]$, define the Boolean moments

$$b_{j,S}^b = \sum_{e \in \{0,1\}^k} b_{j,e} e_S.$$

Summing each coordinate of the Boolean cube separately gives

$$b_{j,S}^b = (-1)^k \sum_{\substack{T \subseteq [k] \\ T \cup S = [k]}} A_{j,T}. \quad (\text{II.146.3})$$

In particular, for an array of degree t ,

$$b_{j,S}^b = 0 \quad (|S| < k - t), \quad (\text{II.146.4})$$

$$b_{j,S}^b = (-1)^k A_{j,[k] \setminus S} \quad (|S| = k - t). \quad (\text{II.146.5})$$

For each $A \subseteq [k+1]$, the associated feature relation is

$$\sum_{j,e} b_{j,e} (v_{j,e})_A = 0. \quad (\text{II.146.6})$$

There are 2^{k+1} such relations, including $A = \emptyset$. The universal line

$$\mathcal{U} = \{\lambda(\sigma_j \pi(e))_{j,e} : \lambda \in \mathbb{F}_q\} \quad (\text{II.146.7})$$

satisfies all of them for every arrangement. Indeed, the alternating sum on a cube vanishes unless a squarefree horizontal feature contains every horizontal coordinate. In that last case the two remaining sums are multiples of $\sum_j \sigma_j = 0$ and $\sum_j \sigma_j z_j = 0$.

In estimating nonuniversal certificates, restrict first to arrangements with every $h_i \neq 0$. This restriction is justified later because a zero side automatically makes (II.145.1) hold for every map.

Lemma II.146.2 (Necessary top-degree balances). *For $0 \leq t \leq k$, let $\mathcal{V}_t$ be the space of arrays whose polynomials in (II.146.2) have degree at most t , and which satisfy*

$$\sum_{j=1}^m A_{j,T} = 0 \quad (|T| = t). \quad (\text{II.146.8})$$

Then $\dim \mathcal{V}_t = D_t$. Every degree- t array satisfying all feature relations on an arrangement with nonzero sides belongs to $\mathcal{V}_t$.

Proof. There are mc_t independent interpolation coefficients of degree at most t . For each of the $\binom{k}{t}$ top monomials, (II.146.8) is one independent equation, giving exactly the asserted dimension.

For a horizontal coordinate set $A \subseteq [k]$, expansion at the anchor gives

$$\sum_e b_{j,e}(x_j + e \odot h)_A = \sum_{S \subseteq A} h_S b_{j,S}^b(x_j)_{A \setminus S}. \quad (\text{II.146.9})$$

Choose $|A| = k - t$. All terms except $S = A$ vanish by (II.146.4). The feature relation is therefore

$$0 = h_A \sum_j b_{j,A}^b = (-1)^k h_A \sum_j A_{j,[k] \setminus A}.$$

Because $h_A \neq 0$, this is (II.146.8). $\square$

The spaces $\mathcal{V}_t$ need not be nested. We never subtract their upper bounds to estimate a stratum. The coefficient arrays are first partitioned by their exact degree, and each resulting mass is bounded above by the mass of the appropriate $\mathcal{V}_t$.

II.146.3 A rank estimate valid for arbitrary field coefficients

The next estimate extends the degree-sensitive certificate argument of [161] from its small coefficient alphabet to arbitrary field coefficients. That extension is essential for a Fejér kernel of arbitrary bandwidth.

Lemma II.146.3 (Probability of a nonuniversal certificate). *Fix $b \in \mathbb{F}_q^s \setminus \mathcal{U}$ of parity-removed degree t . For every fixed $h \in (\mathbb{F}_q^\times)^k$, choose the anchors uniformly in $(\mathbb{F}_q^k)^m$ and the levels uniformly in W . The probability that all feature relations (II.146.6) hold is at most $Q_t(q)$. Consequently the number of successful arrangements with nonzero sides is at most $q^M Q_t(q)$.*

Proof. First let $t = 0$. Then $b_{j,e} = a_j \pi(e)$ for some nonzero vector a . The full horizontal feature requires $\sum_j a_j = 0$, and adjoining the level coordinate requires $a \cdot z = 0$. If the first condition fails there are no successful arrangements. Otherwise a is not a multiple of σ , since the universal line was excluded. The linear form $a \cdot z$ is therefore nonzero on $W = \sigma^\perp$ and vanishes on exactly a fraction q^{-1} of W . This proves the degree-zero bound.

Now suppose $t \geq 1$. Choose $S_0 \subseteq [k]$, of size $k - t$, for which

$$b^b = (b_{1,S_0}^b, \dots, b_{m,S_0}^b) \neq 0.$$

Such a set exists by (II.146.5). For brevity write $b_j^b = b_{j,S_0}^b$. The lowest horizontal and vertical feature relations give

$$\sum_j b_j^b = 0, \quad \sum_j b_j^b z_j = 0. \quad (\text{II.146.10})$$

For each $i \notin S_0$, the feature with horizontal set $S_0 \cup \{i\}$ gives an affine linear equation in the anchors. Its terms involving $x_{j,i}$ are

$$h_{S_0} \sum_j b_j^b x_{j,i}.$$

All remaining anchor terms use coordinates $x_{j,a}$ with $a \in S_0$; condition on those coordinates. In particular, choosing one j_0 with $b_{j_0}^b \neq 0$ gives a diagonal nonsingular minor in the t variables $x_{j_0,i}$,

$i \notin S_0$. For every choice of the other variables, at most a fraction q^{-t} of the anchors can pass these tests.

If $b^\flat$ is not proportional to σ , the second condition in (II.146.10) independently holds on a fraction q^{-1} of W . Thus the success probability is at most q^{-t-1} .

Suppose instead $b^\flat = c\sigma$, with $c \neq 0$. When the level vector z is nonconstant, choose a, b with $z_a \neq z_b$. For each $i \notin S_0$, use both the horizontal test just described and the same test multiplied by the level coordinate. On the variables $x_{a,i}, x_{b,i}$ their coefficient matrix is

$$ch_{S_0} \begin{pmatrix} \sigma_a & \sigma_b \\ z_a \sigma_a & z_b \sigma_b \end{pmatrix}.$$

Its determinant is nonzero. The t blocks use disjoint variables, so their combined rank is $2t$. For constant z , the rank- t argument still applies. The constant vectors form a one-dimensional subspace of W , which has dimension $m - 1$. The fraction of constant level vectors is consequently $q^{-(m-2)}$, and the total probability is at most

$$(1 - q^{-(m-2)})q^{-2t} + q^{-(m-2+t)}. \quad (\text{II.146.11})$$

At $t = 1$, this is exactly $q^{-2} + (q - 1)q^{-m} = Q_1(q)$. For $t \geq 2$, after division by q^{-t-1} it is at most $q^{1-t} + q^{3-m} \leq 2/q < 1$, since $p > 2$. This proves every stated probability. There are fewer than q^k choices of nonzero h , yielding the final counting bound. $\square$

The exception at degree one is thus geometrically explicit: all levels coincide, and the extra vertical rank disappears. Suppressing this event would incorrectly replace Q_1 by q^{-2} . The exhaustive small-field check in Section V.132 attains the stated Q_1 .

II.147 The random restriction and its collision correction

Source 35, delivered Section 4.

II.147.1 A finite positive Fourier selector

For $|z| = 1$, use the normalized Fejér kernel

$$K_L(z) = \left| \frac{1}{L} \sum_{a=0}^{L-1} z^a \right|^2 = \sum_{n=-(L-1)}^{L-1} c_n z^n, \quad c_n = \frac{L - |n|}{L^2}. \quad (\text{II.147.1})$$

Thus $0 \leq K_L \leq 1$, every $c_n \geq 0$, $\sum_n c_n = 1$, and

$$\mu := c_0 = L^{-1} = \max_n c_n.$$

Because $p > 2(L - 1)$, the support integers n embed injectively in $\mathbb{F}_p \subseteq \mathbb{F}_q$. We regard c_n as a probability law on that field subset.

Choose a nontrivial additive character $\chi : \mathbb{F}_q \rightarrow \mathbb{C}$. Let $v \in \mathbb{F}_q$ and $u_A \in \mathbb{F}_q$, for every $A \subseteq [k + 1]$, be independent uniform variables. Conditional on these seeds, include each *distinct physical point* $y \in B$ independently with probability

$$K_L \left(\chi \left(v\phi(y) + \sum_{A \subseteq [k+1]} u_A y_A \right) \right). \quad (\text{II.147.2})$$

Call the resulting set B' . The empty feature is included. By orthogonality, its random coefficient makes the unconditional inclusion probability of each physical point exactly μ . All random choices here have finite support; no continuous phase averaging is required.

For a fixed arrangement, first multiply the retention probabilities over its s labels, whether or not some labels coincide. Fourier expansion and averaging the seeds give

$$\begin{aligned} & \mathbb{E}_{\text{seeds}} \prod_{j,e} K_L \left(\chi \left(v\phi(v_{j,e}) + \sum_A u_A(v_{j,e})_A \right) \right) \\ &= \sum_{b \in \mathbb{F}_q^s} \left(\prod_{j,e} c_{b_{j,e}} \right) \mathbf{1}_{\{\text{all feature relations hold}\}} \mathbf{1}_{\{\sum_{j,e} b_{j,e} \phi(v_{j,e}) = 0\}}. \end{aligned} \quad (\text{II.147.3})$$

Coefficients outside the embedded support have weight zero. Every summand is nonnegative. This positivity allows us to retain a useful signal and bound each unwanted family separately.

Lemma II.147.1 (The universal Fourier signal). *On a respected arrangement, the contribution of $\mathcal{U}$ to (II.147.3) is*

$$S_s(L) = \sum_{n=-(L-1)}^{L-1} c_n^s = \mu^s R_s(L). \quad (\text{II.147.4})$$

On an unrespected arrangement its contribution is exactly μ^s .

Proof. Every array in $\mathcal{U}$ satisfies the feature relations. An array from this line with positive weight must have the form $b_{j,e} = n\sigma_j\pi(e)$ with one integer $|n| < L$: this follows from the injectivity of the support in the field and its symmetry. Its product weight is c_n^s . For a respected arrangement it satisfies the target relation for every such n . For an unrespected arrangement the target relation is n times a nonzero element of $\mathbb{F}_q$, so only $n = 0$ can contribute. This proves both statements. $\square$

II.147.2 Why physical coincidences require a separate term

If labels coincide, the survival probability of the arrangement is the product over its *distinct* points, whereas (II.147.3) used every label. Because each factor lies in $[0, 1]$, actual survival is at least the labelled product. This is favorable for the respected count. For the unrespected count it requires an upper correction.

Lemma II.147.2 (Coincident vertices of an unrespected arrangement). *An unrespected arrangement has all $h_i \neq 0$. The total number of arrangement parameters with nonzero sides and at least one coincident pair of vertex labels is at most*

$$C_{\text{coll}} q^{M-k-1}. \quad (\text{II.147.5})$$

For any such arrangement in B , its unconditional actual survival probability is at most μ .

Proof. If $h_i = 0$, pair the terms e and $e + \mathbf{e}_i$ in each cube of (II.145.1). Their physical vertices are equal and their signs are opposite, so the arrangement is respected by every map. This proves the first assertion.

With nonzero sides, distinct labels in the same cube are different points. A coincidence between cubes $j \neq \ell$ requires

$$z_j = z_\ell, \quad x_\ell - x_j = \delta \odot h \quad \text{for some } \delta \in \{-1, 0, 1\}^k. \quad (\text{II.147.6})$$

For a fixed unordered pair j, ℓ and δ , these are $k+1$ independent linear conditions on the arrangement parameter space. In particular the form $z_j - z_\ell$ is independent of $\sigma \cdot z$, because $m \geq 4$ and σ has nonzero coordinates outside j, ℓ . The locus has exactly q^{M-k-1} parameters before the nonzero-side restriction. There are $\binom{m}{2} 3^k$ loci, so the union bound proves (II.147.5).

Survival implies that any one specified vertex survives. That vertex has unconditional inclusion probability μ , which proves the last assertion. $\square$

Counting raw pairs of Boolean labels would introduce 4^k in this estimate. Grouping pairs that define the same δ produces 3^k . More importantly, the compensation is μ per possible collision, rather than one. The empty feature is what gives this exact marginal probability.

II.147.3 Completion of the selection proof

Proof of Theorem II.145.2. Let $X = G_\phi(B')$ and $Y = T(B') - G_\phi(B')$. For every respected arrangement, actual survival is at least its labelled product. By the nonnegative Fourier expansion and Lemma II.147.1,

$$\mathbb{E}X \geq S_s(L)G_\phi(B) \geq \alpha\mu^s R_s(L)T_0. \quad (\text{II.147.7})$$

For unrespected arrangements, first bound the expected labelled product. Their number is at most $T(B) - G_\phi(B) \leq (1 - \alpha)T_0$, and the universal line contributes μ^s to each. For a nonuniversal array of exact degree t , successful nonzero-side arrangements are possible only if the array is in $\mathcal{V}_t$, by Lemma II.146.2. The product-law mass of that space is at most μ^{s-D_t} , by Lemma II.146.1. Each fixed array succeeds on at most $q^M Q_t(q)$ arrangements by Lemma II.146.3. We may discard both membership in B and the target relation when taking an upper bound. The total expected labelled contribution of these arrays is therefore at most

$$q^M \sum_{t=0}^k \mu^{s-D_t} Q_t(q).$$

The universal line has already been removed, so the degree-zero probability bound is being used only on its valid domain.

For arrangements with distinct labels, actual survival is the labelled product. For the remaining unrespected arrangements, add their actual survival probabilities as a compensation. This may overcount their labelled products, which is harmless. Lemma II.147.2 bounds the compensation by $C_{\text{coll}}\mu q^{M-k-1}$. Thus

$$\begin{aligned} \mathbb{E}Y &\leq \mu^s(1 - \alpha)T_0 \\ &\quad + q^M \left(\sum_{t=0}^k \mu^{s-D_t} Q_t(q) + C_{\text{coll}}\mu q^{M-k-1} \right). \end{aligned} \quad (\text{II.147.8})$$

Subtracting (II.147.8) from η times (II.147.7), and using $q^M = \beta^{-(m-1)}T_0$, gives

$$\mathbb{E}(\eta X - Y) \geq \mu^s T_0 F_*.$$

Some outcome consequently satisfies $\eta X - Y \geq \mu^s T_0 F_* > 0$. For that outcome,

$$X \geq \frac{\mu^s T_0 F_*}{\eta}, \quad Y \leq \eta X,$$

proving the count and purity assertions. Finally the general bound $T(B') \leq (|B'|/q^{k+1})^{m-1}q^M$, together with the lower bound on X , proves the density assertion. $\square$

Remark II.147.3 (A larger class of target groups). The same theorem holds for a map into $\mathbb{F}_q^a$, for any positive integer a : use a random linear functional on the target in the selector. More generally, it holds for a finite abelian target H whenever multiplication by every integer $1 \leq n < L$ is injective on H . For example, it suffices that every prime divisor of $|H|$ exceed $L - 1$. Replace the phase in (II.147.2) by

$$\psi(\phi(y)) \chi \left(\sum_A u_A y_A \right),$$

where ψ is a uniform character of H , and apply K_L to this unit complex number. Orthogonality on H now enforces the target relation. The injectivity assumption ensures that no nonzero integer on the universal line annihilates the defect of an unrespected arrangement. Every other part of the proof is unchanged. The field characteristic condition is still required for the coefficient support used by the features.

II.148 The exponent, a fair comparison, and a downstream substitution

Source 35, delivered Section 5.

II.148.1 A weighted comparison without the full degree stratification

To isolate the contribution of the degree calculation, one should compare it with a weighted argument that uses the same feature family but remembers only whether an array is pure parity. This intermediate comparison is useful even though it is weaker than the main theorem.

Proposition II.148.1 (A weighted unstratified bound). *Theorem II.145.2 remains valid if its budget is replaced by*

$$\mathcal{K}_{\text{un}}(L, q) = L^{m-1} q^{-1} + 2L^{s-1} q^{-2} + C_{\text{coll}} L^{s-1} q^{-k-1}. \quad (\text{II.148.1})$$

For fixed k, m, β , this yields the sufficient bandwidth exponent $(s - 1)/2$.

Proof. The nonuniversal pure-parity arrays lie in the balanced $(m - 1)$ -dimensional space $\mathcal{V}_0$, of mass at most μ^{s-m+1} , and each has success probability at most q^{-1} . Every other successful array satisfies the empty-feature equation $\sum_{j,e} b_{j,e} = 0$. This is a nonzero linear equation on $\mathbb{F}_q^s$, so all these arrays together have mass at most μ . Their individual success probabilities are at most $2q^{-2}$, by Lemma II.146.3. The collision correction is unchanged. Normalizing these contributions by $\mu^s q^M$ gives (II.148.1).

For example, it is enough to impose

$$q \geq \max \left\{ 4\beta^{-(m-1)} L^{m-1}, 2\sqrt{(2 + C_{\text{coll}})\beta^{-(m-1)}} L^{(s-1)/2} \right\}. \quad (\text{II.148.2})$$

Then the first term of the budget is at most $\beta^{m-1}/4$; the other two together are at most the same amount, using $q^{-k-1} \leq q^{-2}$. Since $s = m2^k \geq 2m$, the governing power of L is $(s - 1)/2$. $\square$

Thus the advance is more than a substitution of a probability weight for an unweighted union bound. Degree stratification further reduces the exponent from $(s - 1)/2$ to $E_{k,m}$, strictly so when $k \geq 2$. The unstratified proposition above is itself a consequence of the full feature system and the certificate estimates; it is not attributed to the earlier one-monomial report.

II.148.2 The exponent profile

Proposition II.148.2 (Unimodality and asymptotic size). *Let*

$$E_{k,m} = \max_{0 \leq t \leq k} \frac{D_t}{t+1}.$$

The sequence $D_t/(t+1)$ is unimodal and has at most two adjacent maximizing indices. It increases strictly whenever $t < k/2$. Moreover,

$$E_{k,m} \leq \frac{s-1}{2}, \quad \text{strictly if } k \geq 2, \quad E_{k,m} \sim \frac{2s}{k} \quad (k \rightarrow \infty) \quad (\text{II.148.3})$$

for each fixed $m \geq 4$.

Proof. Put $b_t = \binom{k}{t}$, $g_t = D_t/(t+1)$. For $0 \leq t < k$, the binomial recurrence gives the exact identity

$$g_{t+1} - g_t = \frac{mb_t}{(t+1)(t+2)} \left(\frac{(m-1)k - (m-2)t + 2}{m} - \frac{c_t}{b_t} \right). \quad (\text{II.148.4})$$

The first expression in parentheses decreases strictly with t . The second increases strictly, as is seen from

$$\frac{c_t}{b_t} = \sum_{j=0}^t \prod_{i=0}^{j-1} \frac{t-i}{k-t+i+1}.$$

When t increases, each existing nonempty product increases and a positive summand is added. Hence the sign in (II.148.4) changes at most once, from positive to negative, and at most one difference can be zero. If $t < k/2$, then $c_t/b_t \leq t+1$, while

$$\frac{(m-1)k - (m-2)t + 2}{m} - (t+1) = \frac{(m-1)(k-2t) - (m-2)}{m} \geq \frac{1}{m}.$$

This proves the assertions about the profile.

For the elementary upper bound, the $t=0$ term is $m-1 < (s-1)/2$. For $t \geq 1$, use $D_t \leq s-1$ and $t+1 \geq 2$. If $k \geq 2$, the degree-one numerator is strictly below $s-1$, and the terms $t \geq 2$ have denominator at least three, proving strictness.

For the asymptotic, let X be binomial with parameters $k, 1/2$. Then

$$\frac{D_t}{s} = \Pr(X \leq t) - \frac{1}{m} \Pr(X = t).$$

Take $t = \lceil k/2 + k^{2/3} \rceil$. Chebyshev's inequality shows that both $1 - \Pr(X \leq t)$ and $\Pr(X = t)$ tend to zero, while $k/(t+1) \rightarrow 2$. Thus $\liminf kE_{k,m}/s \geq 2$.

For the reverse inequality fix $0 < \delta < 1/2$. If $t \geq (1/2 - \delta)k$, then

$$\frac{kD_t}{s(t+1)} \leq \frac{1}{1/2 - \delta}.$$

If $\sqrt{k} \leq t < (1/2 - \delta)k$, Chebyshev gives $\Pr(X \leq t) \leq 1/(4\delta^2 k)$, so this same ratio is at most $1/(4\delta^2(t+1))$, which tends uniformly to zero. For $t < \sqrt{k}$,

$$c_t \leq (\sqrt{k} + 1)k^{\sqrt{k}}, \quad \frac{kD_t}{s(t+1)} \leq \frac{kc_t}{2^k} \rightarrow 0.$$

Hence the limsup is at most $1/(1/2 - \delta)$. Letting δ decrease to zero proves the asymptotic. $\square$

The maximizing index can be found without comparing large rational numbers: locate the first t for which the parenthesized expression in (II.148.4) is nonpositive. Equality gives exactly two adjacent maximizers. If no such index occurs before k , the maximum is attained at $t = k$.

II.148.3 Eight-arrangements: explicit numbers

In the interface of the original proof, $d = 8$, so $m = 16$, $s = 2^{k+4}$, and $M = 17k + 15$. The original Lemma 15.5 [158] uses

$$J = s2^{s-1} = 2^{2^{k+4}+k+3}$$

and a sufficient condition

$$(\alpha\eta/4)^J \beta^{15} q \geq k 3^s.$$

Its guaranteed retained count is $(\alpha\eta/4)^J \beta^{15} q^{17k+15}$. The recent one-stage Fejér report [160] already replaces this count loss by

$$\frac{\alpha^s \eta^{s-1}}{3^{s-1}(s+1)^s} \beta^{15} q^{17k+15},$$

under its explicit sufficient size condition

$$q \geq (k+1)\beta^{-15} \left(\frac{18(s+1)^2}{(\alpha\eta)^2} \right)^s$$

and its separate characteristic condition. We retain the same powers $\alpha^s \eta^{s-1}$ and replace the field-size exponent $2s$ in that convenient condition by the following smaller values.

k	s	Maximizing t	Earlier $2s$	Weighted $(s-1)/2$	Graded $E_{k,16}$
1	32	1	64	31/2	31/2
2	64	1	128	63/2	23
3	128	2	256	127/2	109/3
4	256	3	512	255/2	59
5	512	3	1024	511/2	203/2
6	1024	4	2048	1023/2	897/5
7	2048	5	4096	2047/2	633/2
8	4096	5	8192	4095/2	1724/3

Table II.148.1: Sufficient inverse- $\alpha\eta$ exponents for eight-arrangements, with k, m, β fixed. The weighted column is Proposition II.148.1, proved here. Constants and density powers are given separately in the theorems.

For $k = 2, m = 16$, the exact dimensions are

$$(D_0, D_1, D_2) = (15, 46, 63), \quad C_{\text{coll}} = 1080.$$

Thus one may take $L = \lceil 65/(\alpha\eta) \rceil$ and require

$$q \geq \max \left\{ 6\beta^{-15} L^{15}, (12\beta^{-15} L^{46})^{1/2}, (6486\beta^{-15} L^{63})^{1/3} \right\}. \quad (\text{II.148.5})$$

The corresponding respected count is at least

$$\frac{\alpha^{64} \eta^{63}}{2 \cdot 66^{64}} \beta^{15} q^{49},$$

with purity at least $1/(1+\eta)$. In particular, the largest bandwidth power in (II.148.5) is 23, not 128. When the density β also varies, the full maximum in (II.148.5) must be retained; the single exponent comparison does not discard those density dependences.

II.148.4 One explicit substitution into the original proof

The preceding improvements can be passed through one adjacent interface without redoing the entire inverse argument. The arrangement-count input supplied by the product-property lemma in the source transcription (Lemma 14.8) is

$$G_\phi(B) \geq \beta^{7 \cdot 4^{k+1}} \gamma^{21k4^{k+1}} q^{17k+15}. \quad (\text{II.148.6})$$

Here $0 < \beta, \gamma \leq 1$, and β is the density of the relevant domain B . We state the consequence conditionally on this numerical input, so that no unstated variant of the product-property definition is needed.

Corollary II.148.3 (A replacement at the Lemma 15.6 interface). *Suppose the eight-arrangement count (II.148.6) holds, and put*

$$A_k = 7 \cdot 4^{k+1} - 15, \quad B_k = 21k4^{k+1}, \quad s = 2^{k+4}.$$

Take

$$L = \left\lceil 2^{44}(s+1)\beta^{-A_k}\gamma^{-B_k} \right\rceil.$$

If (II.145.12) holds with $m = 16$ and this bandwidth, and $p > 2(L-1)$, there is a subset with purity at least $1 - 2^{-44}$ and at least

$$\frac{2^{-44(s-1)-1}}{(s+2)^s} \beta^{sA_k+15} \gamma^{sB_k} q^{17k+15} \quad (\text{II.148.7})$$

respected arrangements.

Proof. In Corollary II.145.3, substitute $\alpha = \beta^{A_k}\gamma^{B_k}$ and $\eta = 2^{-44}$. Equation (II.148.6) is precisely the required lower bound αT_0 . Expanding the parameter powers in (II.145.13) gives (II.148.7). $\square$

[write] **6 October 2026; sources 34–39.** The input (II.148.6) has exactly the corrected exponents of the formal `lemma_14_8` (Sections14_15.lean, line 292; companion `lemma_14_8_holds`, `Proofs14ProductToArrangements.lean`, line 18, at HEAD 89c1cbc38): $\beta^{7 \cdot 4^{k+1}} \gamma^{21k4^{k+1}} N^{17k+15}$, the exponent $21k4^{k+1}$ in place of the printed $7k4^{k+1}$, as the formal `lemma_15_6` records.

This is a proved local substitution with all exponents visible. Obtaining a global bound for $r_\ell(N)$ would additionally require propagating the new largeness hypotheses through every later restriction, modeling step, and density-increment iteration. That separate optimization has not been carried out here.

II.149 Source 38’s Part II: sharp odd-order stability at the Proposition 6.1 endpoint [write]

Added in the seventh write, 6 October 2026.

The next section is source 38’s Section 6, its delivered Part II. *Logically it continues source 31’s chapter on the Proposition 6.1 endpoint* (Section II.104), which gives the exact profile over all finite abelian groups and leaves the odd-order coefficients open; it is printed here because no number printed earlier may move. Source 38’s introduction is Section F.43.

What is proved. For unit vectors v on groups of odd order with $\|v\|_{U^3}^8 \geq 1 - \varepsilon$, the optimal distance $\mathfrak{S}_{\text{odd}}(\varepsilon)$ from quadratic coset phases is $\varepsilon/4 + 3\varepsilon^2/16 + (\sqrt{2}/64)\varepsilon^{5/2} + O(\varepsilon^3)$, uniformly

in the group; the fractional term comes from characters of order three and disappears when $\gcd(|G|, 6) = 1$ (Theorem II.150.11); through a uniform preliminary stability (Theorem II.150.1, credited to source 31), an explicit local envelope (Theorem II.150.5) and tangent rigidity at order-three characters. Over all groups source 31's profile has the second coefficient $5/16$, attained on $\mathbb{Z}/2\mathbb{Z}$ (note after the theorem). This answers source 31's odd-order question (item II.46).

Relation to the formal statements (checked at HEAD 89c1cbc38). `proposition_6_1` (Section 06_07.lean, line 37) has the exact companion `proposition_6_1_holds` (Proofs05_10.lean, line 505); no near-extremal stability statement is formal. Nothing in this chapter is formalized.

II.150 The sharp odd-order local stability coefficients

Source 38, delivered Section 6.

Delivered part heading of source 38: Sharp odd-order stability for the critical uniformity inequality

The preceding project report proves an odd-order envelope with coefficient 12 in front of $(1-t)^2 t^2$ and leaves its optimal local coefficient open [172]. We determine that coefficient and the next, fractional-power term. The latter distinguishes groups with three-torsion from groups whose order is coprime to six. The resulting statement is a sharp asymptotic modulus, with an explicit finite remainder.

The classification of critical-norm extremizers and qualitative near-extremizers is established by Eisner–Tao [7, Theorem 1.4]. We begin with the finite argument and the explicit preliminary estimate needed here, including their proofs. This makes the quantitative conclusions independent of any unproved assertion in a preceding report. The original Fourier-to-structure setting is Gowers' Proposition 6.1 [1]; the Weyl representation also connects the calculation to stabilizer-state testing [125].

II.150.1 Normalization and quadratic coset phases

Let G be a finite abelian group, $N = |G|$, and let $\langle u, v \rangle = \mathbb{E}_x u(x) \overline{v(x)}$ be the inner product, linear in its first argument. Write $\mathbb{T} = \{z \in \mathbb{C} : |z| = 1\}$. A phase $Q : K \rightarrow \mathbb{T}$ on a subgroup $K \leq G$ is *quadratic* if all its third multiplicative derivatives equal one. Equivalently, for every $h \in K$,

$$Q(x) \overline{Q(x-h)} = c_h \chi_h(x), \quad |c_h| = 1, \quad \chi_h \in \widehat{K}.$$

To see the equivalence, a unit phase with constant first differences is a constant times a character; apply this observation to the first derivatives of Q . A normalized quadratic coset phase is a function

$$\psi(x) = \sqrt{\frac{N}{|K|}} \mathbf{1}_{x_0+K}(x) Q(x-x_0). \quad (\text{II.150.1})$$

Let $\mathcal{Q}_{\text{cos}}(G)$ be this family, including arbitrary constant unit phases. For $\|v\|_2 = 1$, put

$$\kappa(v) = \|v\|_{U^3}^8, \quad \mathcal{F}(v) = \max_{\psi \in \mathcal{Q}_{\text{cos}}(G)} |\langle v, \psi \rangle|^2.$$

The maximum exists, as follows also from the finite eigenline classification below. For a nonzero unnormalized function f , take $v = f/\|f\|_2$; then $\kappa(v) = \|f\|_{U^3}^8 / \|f\|_2^8$.

The phase space $V = G \times \widehat{G}$ has size N^2 . For $z = (h, \chi)$ define

$$W_z = M_\chi T_h, \quad (M_\chi u)(x) = \chi(x)u(x), \quad T_h u(x) = u(x - h).$$

These unitary operators multiply according to addition of their labels, up to unit scalars. Their commutator is

$$W_z W_y = B(z, y) W_y W_z, \quad B((h, \chi), (k, \eta)) = \chi(k) \overline{\eta(h)}. \quad (\text{II.150.2})$$

The alternating bicharacter B is nondegenerate, by character separation on G and on $\widehat{G}$.

For $\|v\|_2 = 1$ set $q_v(z) = |\langle v, W_z v \rangle|^2$. Fourier Parseval and the derivative formula for the U^3 norm give

$$0 \leq q_v \leq 1, \quad \sum_{z \in V} q_v(z) = N, \quad \kappa(v) = \frac{1}{N} \sum_{z \in V} q_v(z)^2 \leq 1. \quad (\text{II.150.3})$$

Indeed $\langle v, W_{(h, \chi)} v \rangle$ is the normalized Fourier coefficient of $v(x) \overline{v(x - h)}$ at χ . Its squared sum in χ is $\mathbb{E}_x |v(x)|^2 |v(x - h)|^2$; summing in h gives N . Its fourth-power sum in χ is the fourth power of the derivative's U^2 norm, and averaging in h gives $\kappa(v)$. The inequality follows from $q_v^2 \leq q_v$ and proves the critical U^3 – L^2 bound with coefficient one.

II.150.2 An explicit preliminary stability theorem

We prove the following estimate on all finite abelian groups. The constant and high-coefficient argument are credited to [172]; they serve here as a self-contained prerequisite.

Theorem II.150.1 (Uniform preliminary stability). *If $\|v\|_2 = 1$ and $\varepsilon = 1 - \kappa(v) < 1/120$, then*

$$\mathcal{F}(v) \geq 1 - 20\varepsilon. \quad (\text{II.150.4})$$

The bound is independent of the order and exponent of G .

Lemma II.150.2 (Projective overlap and additive closure). *If $q_v(z), q_v(y) \geq 1 - a$ with $0 \leq a \leq 1/2$, then*

$$q_v(z - y) \geq (1 - 2a)^2.$$

If a nonempty finite subset S of an abelian group satisfies $|S - S| < 3|S|/2$, then $S - S$ is a subgroup.

Proof. Decompose the unit vectors $W_z v, W_y v$ along v and its orthogonal complement. After adjusting their constant phases, their inner product has absolute value at least $(1 - a) - a = 1 - 2a$. The projective Weyl multiplication law identifies its squared modulus with $q_v(z - y)$.

For the second assertion put $D = S - S$. If $d = b - c \in D$, the sets $S - b, S - c \subseteq D$ have intersection of size greater than $|S|/2$. Equivalently more than half of $x \in S$ satisfy $x + d \in S$. For any $d, e \in D$, the two corresponding subsets of S intersect; for an x in their intersection, both $x + d, x + e$ belong to S . Thus $d - e \in D$, so the nonempty set D is closed under subtraction and is a subgroup. $\square$

Call a subgroup $H \leq V$ *isotropic* if $B(h, k) = 1$ for all $h, k \in H$.

Lemma II.150.3 (A gap for noncommuting subgroups). *If $H \leq V$ is not isotropic, then every unit vector v satisfies*

$$\frac{1}{|H|} \sum_{h \in H} |\langle v, W_h v \rangle|^2 \leq \frac{1}{2}.$$

Proof. For an operator A , average its conjugates:

$$\mathcal{P}(A) = \frac{1}{|H|} \sum_{h \in H} W_h A W_h^*.$$

Conjugation removes the projective scalars, so $\mathcal{P}$ is the orthogonal projection, for the Hilbert–Schmidt inner product, onto the operators commuting with all W_h . Put $\rho = |v\rangle\langle v|$ and $\sigma = \mathcal{P}(\rho)$. This is a positive operator of trace one, and

$$\frac{1}{|H|} \sum_{h \in H} |\langle v, W_h v \rangle|^2 = \text{tr}(\rho\sigma) = \text{tr}(\sigma^2).$$

Every eigenspace of σ is invariant under all W_h . None can be one-dimensional: a common eigenline would force every scalar commutator in (II.150.2) to equal one, contrary to the hypothesis. Each positive eigenvalue therefore has multiplicity at least two, and is at most $1/2$. Hence $\text{tr}(\sigma^2) \leq \frac{1}{2} \text{tr} \sigma = 1/2$. $\square$

[write] **6 October 2026; sources 34–39.** This is source 31’s Lemma II.106.4, reproved for completeness; the preliminary stability around it is credited to source 31 by the source. Source 36’s Theorem IV.17.1 is its sharp rank-sensitive form over $\mathbb{F}_p^n$.

Lemma II.150.4 (Isotropic eigenspaces and quadratic coset phases). *If $H \leq V$ is isotropic, then $|H|$ divides N , and each common character eigenspace of its normalized representation has dimension $N/|H|$. In particular, if $|H| = N$, the common eigenvectors form an orthonormal basis indexed by $\hat{H}$, and each is a quadratic coset phase. Conversely, every quadratic coset phase arises this way for an isotropic subgroup of size N .*

Proof. The commuting unitary operators W_h are simultaneously diagonalizable. Choose a common eigenvector with eigenvalues λ_h and put $U_h = \lambda_h^{-1} W_h$. Evaluating the projective product on that eigenvector shows that $U_h U_k = U_{h+k}$. For each $\gamma \in \hat{H}$ its character projector is

$$P_\gamma = \frac{1}{|H|} \sum_{h \in H} \overline{\gamma(h)} U_h.$$

Every nonidentity Weyl operator has trace zero: a nonzero translation has no diagonal entries, while a nontrivial pure modulation has character sum zero. Thus $\text{tr} P_\gamma = N/|H|$. This positive number is an integer, proving the divisibility and dimension statements.

Assume now that $|H| = N$. Let K be its projection into G and let $D = \{\chi : (0, \chi) \in H\}$. Isotropy gives $D \subseteq K^\perp$; the equality $|H| = |K||D| = N$ forces $D = K^\perp$. The modulation eigenrelations restrict the support of a common eigenvector to one coset $x_0 + K$. The translation eigenrelations for $h \in K$ force that entire coset into the support, with constant modulus. Normalization fixes the modulus to $\sqrt{N/|K|}$. The same relations state that each derivative of its phase on K is a constant times a character, which is precisely quadraticity.

Conversely consider (II.150.1). For every $h \in K$ the derivative character on K has $|K^\perp|$ extensions to G . Each extension gives a Weyl operator preserving the line through ψ . The labels of all line-preserving Weyl operators form a subgroup, by the projective multiplication law, and its commutators must be one since they act on a nonzero common eigenvector. No label whose translation lies outside K can preserve that line. Consequently this isotropic subgroup has exactly $|K||K^\perp| = N$ elements. The finite-group character-extension fact used here follows, for example,

by extending a character one cyclic generator at a time and choosing a suitable root of its already specified value.

There are finitely many subgroups of V and exactly N eigenlines for each isotropic subgroup of size N . The classification therefore also proves that $\mathcal{Q}_{\text{cos}}(G)$ is finite modulo constant unit phases, and that the maximum defining $\mathcal{F}(v)$ is attained. $\square$

Proof of Theorem II.150.1. Write $q = q_v$ and put $S = \{z : q(z) \geq 19/20\}$. Equation (II.150.3) gives $\sum_z q(z)(1 - q(z)) = N\varepsilon$, so the q -mass outside S is at most $20N\varepsilon$. Hence

$$|S| \geq \sum_{z \in S} q(z) \geq (1 - 20\varepsilon)N.$$

The overlap lemma puts $S - S$ inside $\{q \geq 4/5\}$, so $|S - S| \leq 5N/4$. Since $\varepsilon < 1/120$,

$$|S - S| \leq 5N/4 < \frac{3}{2}(1 - 20\varepsilon)N \leq 3|S|/2.$$

Thus $H = S - S$ is a subgroup. Every coefficient there is at least $4/5$, so Lemma II.150.3 makes H isotropic. Also $|H| \geq |S| > 5N/6 > N/2$, and divisibility in Lemma II.150.4 implies $|H| = N$.

Since $q(0) = 1$, one has $0 \in S$ and $S \subseteq H$. Expand v in the common eigenbasis of H , with coefficient probabilities p_γ . Character orthogonality gives

$$\sum_\gamma p_\gamma^2 = \frac{1}{N} \sum_{h \in H} q(h) \geq 1 - 20\varepsilon.$$

The largest p_γ is at least their sum of squares because their sum is one. Its eigenvector is a quadratic coset phase, proving (II.150.4). $\square$

At $\varepsilon = 0$, the argument and the mass identity also recover the exact extremizers: they are precisely the normalized quadratic coset phases. The cutoff $1/120$ and preliminary coefficient 20 are sufficient constants; their optimality is not asserted.

II.150.3 The explicit odd-order local envelope

Theorem II.150.5 (Explicit odd-order local envelope). *Suppose $|G|$ is odd and $\|v\|_2 = 1$. Let ψ be any normalized quadratic coset phase with $|\langle v, \psi \rangle|^2 = 1 - t$, where $0 \leq t \leq 1/6$. Set $F = 1 - t$. Then*

$$\kappa(v) \leq F^4 + 6F^2t^2 + 2\sqrt{2}F^{3/2}t^{5/2} + 104Ft^3 + t^4. \quad (\text{II.150.5})$$

If $3 \nmid |G|$, the term $2\sqrt{2}F^{3/2}t^{5/2}$ may be omitted. The constants 6 and $2\sqrt{2}$, in that order of asymptotic comparison, are optimal uniformly over odd-order groups. No optimality is claimed for 104.

Here “in that order” means that the coefficient of F^2t^2 is first minimized; among envelopes with that coefficient equal to 6, the coefficient of $F^{3/2}t^{5/2}$ is then minimized. The theorem does not assert an exact finite- t optimum.

II.150.4 A Weyl-coset expansion

Choose the isotropic subgroup H of size N associated with ψ by Lemma II.150.4, and normalize its representation so that ψ has character label zero. Its common eigenbasis is indexed by $\Gamma = \widehat{H}$. Nondegeneracy of the commutator pairing gives $|H^\perp| = N^2/|H| = N$, hence $H^\perp = H$. The map sending a Weyl label to its commutator character on H therefore identifies V/H with Γ . A Weyl operator from a coset with label γ shifts eigenvector labels by γ , up to unit scalar factors. This follows by commuting that operator past each U_h . These facts hold for every maximal isotropic subgroup; no change to a coordinate subgroup is assumed.

Choose the global phase so that

$$v = \sqrt{F}\psi + \sqrt{t}w, \quad \langle w, \psi \rangle = 0, \quad \|w\|_2 = 1.$$

The case $t = 0$ is immediate and is henceforth suppressed. Let s_γ be the coefficient probabilities of w in the common eigenbasis. Then $s_0 = 0$, $\sum s_\gamma = 1$. Define

$$A = \sum_\gamma s_\gamma^2, \quad B = \sum_\gamma s_\gamma s_{-\gamma}, \quad C = \sum_{\alpha+\beta=\gamma} s_\alpha s_\beta s_\gamma.$$

Since $|\Gamma|$ is odd, its nonzero elements form disjoint pairs $\{\gamma, -\gamma\}$. If m_i are their total probabilities, then

$$A + B = \sum_i m_i^2 \leq 1, \quad B \leq \frac{1}{2}. \quad (\text{II.150.6})$$

In particular $2A + 4B \leq 3$.

For $z \notin H$ put

$$l_z = \langle \psi, W_z w \rangle + \langle w, W_z \psi \rangle, \quad d_z = \langle w, W_z w \rangle.$$

Write $\int_{V \setminus H} a = N^{-1} \sum_{z \notin H} a(z)$; this is a finite measure, not necessarily a probability measure.

Here is an explicit formula for the coset coordinates, including the phase factors. Write $w = \sum_\beta b_\beta \psi_\beta$, where $U_h \psi_\beta = \beta(h) \psi_\beta$ and $b_0 = 0$. Choose a representative z_γ for a Weyl coset. There are unit scalars $u_{\gamma, \beta}$ such that, after removing the same common unit factor from all coefficients at $z_\gamma + h$, the ambiguity coefficient of a vector with coordinates c_β is

$$\sum_{\beta \in \Gamma} c_{\beta+\gamma} \overline{c_\beta} u_{\gamma, \beta} \overline{\beta(h)}. \quad (\text{II.150.7})$$

This follows by first applying U_h and then the label-shifting operator W_{z_γ} . Thus, for $\gamma \neq 0$, the linear term is

$$l_\gamma(h) = b_\gamma u_{\gamma, 0} + \overline{b_{-\gamma}} u_{\gamma, -\gamma} \gamma(h).$$

The same common factor multiplies l and d , and cancels from every product $|l|^4$, $|l|^2 |d|^2$, and $l^2 \bar{l} \bar{d}$ used below. In particular the argument does not require the projective scalar factors to vanish separately.

Lemma II.150.6 (The exact fourth-order coefficient). *One has*

$$\int_{V \setminus H} |l|^4 = 2A + 4B. \quad (\text{II.150.8})$$

Furthermore, with

$$K(w) = 4 \operatorname{Re} \int_{V \setminus H} l^2 \bar{l} \bar{d},$$

there exist real numbers $J(w), M(w)$ such that

$$\begin{aligned} \kappa(v) &= F^4 + 6F^2t^2(A + B) + F^{3/2}t^{5/2}K(w) \\ &\quad + Ft^3J(w) + F^{1/2}t^{7/2}M(w) + t^4\kappa(w), \end{aligned} \quad (\text{II.150.9})$$

$$0 \leq J(w) \leq 4 + 6\sqrt{3}, \quad |M(w)| \leq 43^{1/4}, \quad \kappa(w) \leq 1. \quad (\text{II.150.10})$$

Proof. On H the coefficient of v is, up to a unit scalar, $F + t\hat{s}(h)$. Character orthogonality gives its fourth moment as

$$F^4 + F^2t^2(4A + 2B) + 4Ft^3C + t^4E(s),$$

where $E(s) = \sum_{a+b=c+d} s_a s_b s_c s_d$. The term with one nonzero label vanishes because $s_0 = 0$.

Consider a Weyl coset with nonzero shift label γ . Up to a common unit scalar, the two summands of l are two characters on H with coefficient magnitudes $\sqrt{s_\gamma}$ and $\sqrt{s_{-\gamma}}$. The ratio character has label γ . Both that character and its square are nontrivial since Γ has odd order. Therefore the fourth moment over this coset is

$$s_\gamma^2 + s_{-\gamma}^2 + 4s_\gamma s_{-\gamma}.$$

Summing over nonzero shift labels proves (II.150.8).

Outside H the Weyl coefficient is $\sqrt{Ft}l + td$. Expand its fourth power. Besides the fourth and fifth-order terms already displayed, the coefficients are

$$J = 4C + \int_{V \setminus H} (2|l|^2|d|^2 + 4(\text{Re}(l\bar{d}))^2), \quad M = 4 \text{Re} \int_{V \setminus H} |d|^2 l \bar{d}.$$

The t^4 terms on and off H add up to $t^4\kappa(w)$. The fourth-order coefficients add to $6(A + B)$, proving the exact identity. Since $C \leq 1$, $\|l\|_4 \leq 3^{1/4}$, and $\|d\|_4 \leq \kappa(w)^{1/4} \leq 1$, Hölder's inequality gives (II.150.10). $\square$

For clarity, the critical estimate for w in this proof does not require an inverse theorem. For any vectors u_1, u_2 , the mixed ambiguity coefficient satisfies

$$\left(\frac{1}{N} \sum_z |\langle u_1, W_z u_2 \rangle|^4 \right)^{1/4} \leq \|u_1\|_2 \|u_2\|_2. \quad (\text{II.150.11})$$

Indeed the maximum squared coefficient is at most $\|u_1\|_2^2 \|u_2\|_2^2$, and Weyl orthogonality gives total squared mass $N\|u_1\|_2^2 \|u_2\|_2^2$. Multiplying these two bounds proves (II.150.11). Taking $u_1 = u_2 = w$ gives $\kappa(w) \leq 1$.

II.150.5 The order-three resonance

Lemma II.150.7 (A single opposite pair). *Let w_0 be supported on the two eigenvector labels $\{\gamma, -\gamma\}$, with $\|w_0\|_2 \leq 1$. Then*

$$|K(w_0)| \leq 2\sqrt{2}.$$

If $3\gamma \neq 0$, then $K(w_0) = 0$.

Proof. The linear coefficient $l(w_0)$ is supported on Weyl cosets with shift labels $\pm\gamma$. The coefficient $d(w_0)$, away from H , is supported on shift labels $\pm 2\gamma$. These sets meet only if $3\gamma = 0$, proving the vanishing assertion.

In the remaining case choose the shift- γ coset. Its linear coefficient has the form $b + c\chi$, up to unit scalars, where $|b| = u$, $|c| = v$, and χ is an order-three character. Its quadratic coefficient has the form $d\chi^{-1}$, with $|d| = uv$. In the average of $l^2\bar{l}\bar{d}$, precisely the powers χ^{-1} and χ^2 can contribute after multiplication by χ . Their total absolute value is at most $u^3v^2 + u^2v^3$. The other shift coset has the same bound. Consequently

$$|K(w_0)| \leq 8u^2v^2(u+v) \leq 2\sqrt{2},$$

using $u^2 + v^2 \leq 1$. Equality in the scalar bound occurs at $u = v = 1/\sqrt{2}$. $\square$

Lemma II.150.8 (Stability of the fifth-order coefficient). *Put $D = 1 - A - B$. Then*

$$K(w) \leq K_* + 20 \cdot 3^{3/4} \sqrt{D}, \quad K_* = \begin{cases} 2\sqrt{2}, & 3 \mid N, \\ 0, & 3 \nmid N. \end{cases} \quad (\text{II.150.12})$$

Proof. Equation (II.150.6) implies that the greatest opposite-pair mass is at least $1 - D$. Let w_0 be the unnormalized orthogonal projection of w onto that pair. Then $\|w - w_0\|_2 \leq \sqrt{D}$. The fourth-moment formula for l , by homogeneity, gives

$$\|l(w) - l(w_0)\|_4 \leq 3^{1/4} \sqrt{D}, \quad \|l(w)\|_4, \|l(w_0)\|_4 \leq 3^{1/4}.$$

The mixed ambiguity estimate (II.150.11) gives

$$\|d(w) - d(w_0)\|_4 \leq (\|w\|_2 + \|w_0\|_2) \|w - w_0\|_2 \leq 2\sqrt{D}.$$

Expand the difference of $l^2\bar{l}\bar{d}$ at w and w_0 , using three terms for the cubic factor and one for d . Hölder gives

$$|K(w) - K(w_0)| \leq 4(3 \cdot 3^{3/4} + 2 \cdot 3^{3/4}) \sqrt{D} = 20 \cdot 3^{3/4} \sqrt{D}.$$

Lemma II.150.7 concludes the proof. If $3 \nmid N$, no label has order three because $|\Gamma| = N$. $\square$

Proof of Theorem II.150.5. Insert (II.150.12) into the exact expansion and set $L = 20 \cdot 3^{3/4}$. Completing the square gives

$$-6F^2t^2D + LF^{3/2}t^{5/2}\sqrt{D} \leq \frac{L^2}{24}Ft^3 = 50\sqrt{3}Ft^3.$$

The remaining terms are at most

$$F^4 + 6F^2t^2 + K_*F^{3/2}t^{5/2} + (4 + 56\sqrt{3})Ft^3 + 4 \cdot 3^{1/4}F^{1/2}t^{7/2} + t^4.$$

Since $t/F \leq 1/5$ for $t \leq 1/6$, the last two intermediate remainder coefficients satisfy

$$4 + 56\sqrt{3} + 4 \cdot 3^{1/4} \sqrt{t/F} \leq 4 + 56\sqrt{3} + \frac{4 \cdot 3^{1/4}}{\sqrt{5}} < 104.$$

This proves the explicit envelope. Sharpness of the two asymptotic coefficients follows from the examples below. $\square$

Corollary II.150.9 (Tangent rigidity at the optimal coefficients). *Consider any sequence of unit vectors on finite abelian groups of odd order, and choose a maximizing quadratic coset phase for each vector. After adjusting its constant phase, write*

$$v = \sqrt{F}\psi + \sqrt{t}w, \quad F = \mathcal{F}(v) = 1 - t, \quad t \downarrow 0.$$

If

$$\kappa(v) = F^4 + 6F^2t^2 + 2\sqrt{2}F^{3/2}t^{5/2} + o(t^{5/2}), \quad (\text{II.150.13})$$

then, in the common eigenbasis associated with ψ , the tangent w concentrates in a pair $\{\gamma, -\gamma\}$ with $3\gamma = 0$. More precisely, these labels have order three for all sufficiently late terms, the two coefficient moduli tend to $1/\sqrt{2}$, and the orthogonal projection w_0 to that pair satisfies

$$1 - A - B = o(\sqrt{t}), \quad \|w - w_0\|_2 = o(t^{1/4}).$$

No classification of the two coefficient phases is asserted.

Proof. Put $D = 1 - A - B \geq 0$ and $K_0 = 2\sqrt{2}$. Comparing (II.150.13) with the exact expansion (II.150.9), using its uniform remainder bounds, gives

$$6\sqrt{F} \frac{D}{\sqrt{t}} = K(w) - K_0 + o(1).$$

The bound (II.150.12) implies that $K(w)$ is bounded, so first $D = O(\sqrt{t})$ and hence $D \rightarrow 0$. It then gives $K(w) \leq K_0 + o(1)$. Since the left side of the displayed identity is nonnegative, both $D/\sqrt{t} \rightarrow 0$ and $K(w) \rightarrow K_0$ follow.

Choose a pair of greatest mass and project to it as in Lemma II.150.8. Then $\|w - w_0\|_2^2 \leq D = o(\sqrt{t})$, and the same Lipschitz estimate gives $K(w_0) \rightarrow 2\sqrt{2}$. Lemma II.150.7 forces the pair label to have order three eventually. If its coefficient moduli are a, b , then $a^2 + b^2 \rightarrow 1$ and

$$K(w_0) \leq 8a^2b^2(a + b) \leq 2\sqrt{2}.$$

The limiting upper bound can be attained only when $a^2b^2 \rightarrow 1/4$: use $a + b \leq \sqrt{2}$ and $a^2b^2 \leq 1/4$. Together with $a^2 + b^2 \rightarrow 1$, this implies $a, b \rightarrow 1/\sqrt{2}$, as required. $\square$

II.150.6 An exact three-element sharpness family

Let $e_j = \sqrt{3}\mathbf{1}_{\{j\}}$ be the normalized point-mass basis on $\mathbb{Z}/3\mathbb{Z}$. For $0 \leq t \leq 1/6$ put

$$v_t = \sqrt{1-t}e_0 + \sqrt{t/2}(e_1 + e_2).$$

Then its best quadratic-coset fidelity is $\mathcal{F}(v_t) = 1 - t$. To verify this, note that a quadratic coset phase is either a point mass or has full support and constant modulus one. The latter's squared overlap is at most $(\sqrt{1-t} + \sqrt{2t})^2/3$, which is smaller than $1 - t$ on this interval. (Squaring the equivalent positive inequality reduces the assertion to $6t(1-t) < 1$.)

The norm is exactly

$$\kappa(v_t) = (1-t)^4 + 6(1-t)^2t^2 + 2\sqrt{2}(1-t)^{3/2}t^{5/2} + 4(1-t)t^3 + \frac{1}{2}t^4. \quad (\text{II.150.14})$$

One direct verification expands the cube sum. Of the $3^4 = 81$ affine maps (x, h_1, h_2, h_3) , the numbers with exactly r zero cube vertices are

r	0	2	3	4	8
number	8	32	16	24	1.

The values at every nonzero vertex coincide and are real. Substituting $a = \sqrt{1-t}$ and $b = \sqrt{t/2}$ gives $a^8 + 24a^4b^4 + 16a^3b^5 + 32a^2b^6 + 8b^8$, which is (II.150.14). Here is an elementary derivation of the count. With no nonzero directions there are one map with $r = 8$ and two with $r = 0$. With one nonzero direction there are six direction tuples; each has two translates with $r = 4$ and one with $r = 0$. With two nonzero directions there are twelve direction tuples; each has residue multiplicities 4, 2, 2 on the cube, hence one translate with $r = 4$ and two with $r = 2$. With three nonzero directions there are eight tuples; their residue multiplicities are 3, 3, 2, hence two translates with $r = 3$ and one with $r = 2$. Adding these contributions gives the table. A short exhaustive verification script accompanies the article.

This proves that the coefficient 6 cannot decrease, and that after it is fixed the coefficient $2\sqrt{2}$ cannot decrease. It also disproves the putative remainder-free envelope $\kappa \leq (1-t)^4 + 6(1-t)^2t^2 + t^4$ on all odd-order groups.

For an example with $3 \nmid N$, use $G = \mathbb{Z}/5\mathbb{Z}$, now with normalized point masses $e_j = \sqrt{5} \mathbf{1}_{\{j\}}$, and put $v_t = \sqrt{1-t}e_0 + \sqrt{t}e_1$. Then

$$\mathcal{F}(v_t) = 1 - t, \quad \kappa(v_t) = (1-t)^4 + 6(1-t)^2t^2 + t^4. \quad (\text{II.150.15})$$

Indeed the zero-shift coset contributes $F^4 + 4F^2t^2 + t^4$, and the two nonzero shift cosets ± 1 each contribute F^2t^2 . Full-support phases have squared correlation at most $2/5$, so the point mass is the maximizer for $t \leq 1/6$.

Remark II.150.10 (Finite comparison and the meaning of the coefficient). The earlier envelope is $F^4 + 12F^2t^2 + t^4$ on odd-order groups [172]. Our explicit upper bound is strictly smaller for $0 < t \leq 1/25$, because

$$2\sqrt{2}\sqrt{t/F} + 104t/F \leq \frac{1}{\sqrt{3}} + \frac{13}{3} < 6.$$

More significantly, its two successive asymptotic coefficients are optimal. The exact three-element family shows why there is no remainder-free replacement by the polynomial with coefficient 6. The next contribution is larger than any constant multiple of t^3 as $t \downarrow 0$, and must be retained when three-torsion is allowed.

II.150.7 The optimal local modulus through its fractional term

Define

$$\mathfrak{S}_{\text{odd}}(\varepsilon) = \sup\{1 - \mathcal{F}(v) : |G| \text{ odd}, \|v\|_2 = 1, \kappa(v) \geq 1 - \varepsilon\},$$

where the supremum ranges over all finite abelian groups as well as all such functions. Define $\mathfrak{S}_{6'}$ by restricting to $\gcd(|G|, 6) = 1$.

Theorem II.150.11 (Sharp uniform asymptotic stability). *As $\varepsilon \downarrow 0$,*

$$\mathfrak{S}_{\text{odd}}(\varepsilon) = \frac{\varepsilon}{4} + \frac{3\varepsilon^2}{16} + \frac{\sqrt{2}}{64}\varepsilon^{5/2} + O(\varepsilon^3), \quad (\text{II.150.16})$$

$$\mathfrak{S}_{6'}(\varepsilon) = \frac{\varepsilon}{4} + \frac{3\varepsilon^2}{16} + O(\varepsilon^3). \quad (\text{II.150.17})$$

The implied constants are independent of group order and exponent. For optimal squared distance to a normalized coset phase, replace these right sides respectively by

$$\frac{\varepsilon}{4} + \frac{13\varepsilon^2}{64} + \frac{\sqrt{2}}{64}\varepsilon^{5/2} + O(\varepsilon^3), \quad \frac{\varepsilon}{4} + \frac{13\varepsilon^2}{64} + O(\varepsilon^3).$$

Proof. Theorem II.150.1 gives $1 - \mathcal{F}(v) \leq 20\varepsilon$ when $\varepsilon < 1/120$, on every finite abelian group. Thus $t = 1 - \mathcal{F}(v)$ is at most $1/6$ and is $O(\varepsilon)$ uniformly. Apply Theorem II.150.5 to a maximizing phase and expand:

$$\varepsilon \geq 4t - 12t^2 - K_*t^{5/2} - O(t^3).$$

First this gives $t \leq \varepsilon/4 + O(\varepsilon^2)$; inserting that estimate back into the quadratic and fractional terms gives

$$t \leq \frac{\varepsilon}{4} + \frac{3\varepsilon^2}{16} + \frac{K_*}{128}\varepsilon^{5/2} + O(\varepsilon^3).$$

All constants are uniform because the explicit envelope and preliminary bound are uniform. The exact family (II.150.14) gives the matching lower asymptotic expansion: its deficit is continuous and strictly increasing for all sufficiently small t , with derivative 4 at zero, so every sufficiently small ε is attained. Equation (II.150.15) gives the analogous lower bound when $3 \nmid N$.

Finally optimizing the constant phase of an extremizer gives

$$\min_{\psi \in \mathcal{Q}_{\cos}(G)} \|v - \psi\|_2^2 = 2(1 - \sqrt{\mathcal{F}(v)}) = t + \frac{1}{4}t^2 + O(t^3).$$

This changes $3/16$ to $3/16 + 1/64 = 13/64$ and preserves the $\varepsilon^{5/2}$ coefficient. □

[write] 6 October 2026; sources 34–39. *Answers source 31's odd-order question* (item II.46; checked at the write). Over all finite abelian groups source 31's exact profile $(1 - \sqrt{(1 + \sqrt{1 - 4\varepsilon})/2})/2$ expands as $\varepsilon/4 + 5\varepsilon^2/16 + O(\varepsilon^3)$, with equality on $\mathbb{Z}/2\mathbb{Z}$; on odd-order groups the second coefficient drops to $3/16$, and the term $(\sqrt{2}/64)\varepsilon^{5/2}$ comes from characters of order three, absent when $\gcd(|G|, 6) = 1$. The two profiles are compatible ($3/16 < 5/16$). The local envelope's constant 104 is not claimed optimal.

[write] 7 October 2026; sources 54–59. Source 55's Theorem II.237.1 has the odd-order squared-distance coefficient $(3^d - 1)/(2q^2)$ in every degree, $13/64$ at $d = 3$ as here, for functions bounded by one, where no $\varepsilon^{5/2}$ term occurs: the fractional term above comes from L^2 -normalized near point masses, which are not bounded by one (note after source 55's theorem). Source 55 does not cite this source.

Scope and remaining questions. The critical-norm classification and the preliminary stability estimate are credited prerequisites, proved above for completeness. The new conclusion is the uniform sharp odd-order expansion through $\varepsilon^{5/2}$, together with an explicit finite- t estimate. It refines the local near-extremal regime of the derivative-Fourier framework; no new global inverse theorem is asserted.

Is the exact odd-order local optimum attained by the three-element family (II.150.14) throughout some nonzero interval? What is the sharp coefficient of ε^3 in $\mathfrak{S}_{\text{odd}}$ and $\mathfrak{S}_{6'}$? Does the latter coefficient depend on five-torsion? A sharper remainder retaining the opposite-pair deficit $1 - \sum_i m_i^2$ could help distinguish these possibilities.

II.151 Source 39’s chapter: shared anchors at Lemma 16.10 [write]

Added in the seventh write, 6 October 2026.

The next ten sections are source 39’s Sections 2–11, on the sampling and interpolation step inside Gowers’s Lemma 16.10: one small set of columns serves as interpolation anchors for many fibres at once. *Logically the chapter follows source 19’s sampling sections* (Sections II.67–II.69), *which are its spine*, and source 23’s geometric gluing (Section II.103); it is printed here because no number printed earlier may move. Source 39 re-derives much of source 19 without citing it (note in Section F.44). Its introduction is Section F.44; its repository interface, checks and questions are in Part V.

What is proved, and what is new in this report. Second routes (source 19’s results, re-derived): the common r -subset sampler with an $O(q/\sigma)$ budget, the exact finite minimax as a maximum over class profiles attained by permutation orbits, the dynamic program, conditional-expectation selection, the sharp affine constant $\varphi^3 e^{-\varphi}$ (Theorems II.153.2, II.154.2, II.158.2, Corollary II.156.3). New: a certification rule that keeps sampled targets, with loss at most $tq(m-r)/(m(r+1))$; the closed balanced-and-capped solution of the finite profile problem for that rule, with an explicit quadratic-radical cap for affine classes (Theorems II.155.3, II.155.4; for source 19’s rule the formula fails, note there); the two-regime limit $F_t(\min\{\lambda_t, r/q\})$ (Theorem II.157.3); unisolvent spaces and additive corruption (Corollary II.159.4, Theorem II.159.5); labelled slice covers and label-preserving refinement (Theorems II.160.1, II.161.2). The closing comparison of Lemma 16.10 is not addressed.

Relation to the formal statements (checked at HEAD 89c1cbc38). `lemma_16_10` (Section16.lean, line 851) is open; source 39’s Proposition V.145.1 matches the existing `Section16SampledOrAnchoredOn` interface (`Proofs16ShortColumns.lean`, line 27), a formalization candidate claimed by nobody. Nothing in this chapter is formalized.

[write] **7 October 2026; sources 47–53.** *Lemma 16.10 itself, as catalogued, is false* (source 50, Section II.226): its two packaged premises hold for $F(h, y) = a(y)$, a a balanced word with $4Q$ values, while its conclusion fails; the repository has since formalized the negation (`lemma_16_10_packaged_premises_counterexample`, `Proofs16PackagedPremisesCounterexample.lean`, line 13, at HEAD 06b30ca52), keeps `lemma_16_10` open, and names source 49’s acyclic cover theorem and source 50’s repair interface as candidates for a faithful replacement. The sampling and interpolation of this chapter and of sources 19 and 23 are not affected; they are inputs to any repair.

II.152 A finite model of shared-anchor certification

Source 39, delivered Section 2.

Definition II.152.1 (Row families and ambient mass). Let J be a finite set with $|J| = m \geq 1$, and let H be a finite row set. Assign weights $w_h \geq 0$ with $W = \sum_h w_h > 0$. In each row choose pairwise disjoint sets

$$C_{h,1}, \dots, C_{h,q_h} \subseteq J, \quad B_h = \bigcup_{i=1}^{q_h} C_{h,i}.$$

Empty classes are permitted. The weighted mean class count and domain density are

$$\bar{q} = \frac{1}{W} \sum_h w_h q_h, \quad \delta_B = \frac{1}{mW} \sum_h w_h |B_h|.$$

All mass bounds in this article are normalized by the full ambient mass mW , not by $\sum_h w_h |B_h|$.

Allowing B_h to be a proper subset of J is essential. An intermediate partial graph in the inverse argument need not meet every vertical line in every column. It is also essential for the extremal profiles: in one of the asymptotic regimes it is best to use only part of the column set.

Definition II.152.2 (Certification). Fix an integer $t \geq 1$ and a sample $R \subseteq J$. Define

$$K_h(R) = (B_h \cap R) \cup \bigcup_{\substack{1 \leq i \leq q_h \\ |C_{h,i} \cap R| \geq t}} C_{h,i}. \quad (\text{II.152.1})$$

The normalized uncertified mass is

$$D_{\mathcal{C}}(R) = \frac{1}{mW} \sum_h w_h |B_h \setminus K_h(R)|. \quad (\text{II.152.2})$$

The ambient good set $E_R \subseteq H \times J$ is obtained by deleting just these uncertified points. Thus all points outside the domain B remain in E_R , and all sampled domain points remain in E_R .

For a class C of size s with $a = |C \cap R|$, the discarded count is precisely

$$(s - a) \mathbf{1}_{\{a < t\}}. \quad (\text{II.152.3})$$

There are two contributions to the definition, and neither should be dropped: a class with enough anchors is retained entirely, while a class with too few anchors still retains its sampled points.

Remark II.152.3 (Monotonicity). If $R \subseteq R'$, then $K_h(R) \subseteq K_h(R')$ and $D_{\mathcal{C}}(R') \leq D_{\mathcal{C}}(R)$. Consequently, a result for at most r anchors can be padded to exactly $r \leq m$ distinct anchors without increasing the loss. The statements below use exact-size samples because their probability law and extremal problem are especially clean.

We use the convention $\binom{a}{b} = 0$ when $a \geq 0$ and $b \notin \{0, \dots, a\}$. A hypergeometric variable $\text{Hyp}(m, s, u)$ counts the members of a fixed s -element class in a uniformly chosen u -subset of an m -element set. The case $r = m$ is treated explicitly whenever an $(r + 1)$ -subset appears.

II.153 A cutoff-free shared sampling theorem

Source 39, delivered Section 3.

Lemma II.153.1 (Exact one-class loss). *Let R be a uniformly chosen r -subset of J , where $0 \leq r \leq m$. A class of size s contributes expected normalized loss $\ell_{m,r,t}(s)$ from (F.44.2). If $r < m$ and $Z \sim \text{Hyp}(m, s, r + 1)$, then*

$$\ell_{m,r,t}(s) = \frac{m - r}{m(r + 1)} \mathbb{E}[Z \mathbf{1}_{\{1 \leq Z \leq t\}}]. \quad (\text{II.153.1})$$

For $r = m$, this loss is zero.

Proof. The random variable $A = |C \cap R|$ has probability $\binom{s}{j} \binom{m-s}{r-j} / \binom{m}{r}$ of equalling j . Taking the expectation of (II.152.3) and dividing by m gives (F.44.2).

For the second formula, count ordered pairs (R, x) such that $|R| = r$, $x \in C \setminus R$, and $|C \cap R| < t$. The number of such pairs is

$$\sum_{|R|=r} |C \setminus R| \mathbf{1}_{\{|C \cap R| < t\}}.$$

Map a pair to $T = R \cup \{x\}$, an $(r+1)$ -subset. Conversely, for fixed T , a choice $x \in T \cap C$ gives a valid pair exactly when $1 \leq |T \cap C| \leq t$. Hence the same count is

$$\sum_{|T|=r+1} |T \cap C| \mathbf{1}_{\{1 \leq |T \cap C| \leq t\}}.$$

Divide by $m \binom{m}{r}$ and use $\binom{m}{r+1} / \binom{m}{r} = (m-r)/(r+1)$. When $r = m$, every domain point is sampled, so no point is discarded. $\square$

The shift from an r -sample to an $(r+1)$ -sample is the central combinatorial identity. It converts the size-weighted error into the expectation of a variable bounded by t . No lower bound on s is required.

Theorem II.153.2 (Shared-anchor selection). *For every row family in Definition II.152.1, every $t \geq 1$, and every $0 \leq r \leq m$, there is one set $R \subseteq J$ with $|R| = r$ such that*

$$D_C(R) \leq \frac{1}{W} \sum_h w_h \sum_{i=1}^{q_h} \ell_{m,r,t}(|C_{h,i}|) \quad (\text{II.153.2})$$

$$\leq \min \left\{ \delta_B, \frac{t\bar{q}(m-r)}{m(r+1)} \right\}. \quad (\text{II.153.3})$$

In particular, if every row has at most q classes, then $D_C(R) \leq tq/(r+1)$.

Proof. For a uniform r -subset, linearity of expectation and disjointness of the classes give

$$\mathbb{E}_R D_C(R) = \frac{1}{W} \sum_h w_h \sum_i \ell_{m,r,t}(|C_{h,i}|).$$

At least one sample has loss no greater than this average. Formula (II.153.1) bounds each summand by $t(m-r)/(m(r+1))$ when $r < m$; the endpoint $r = m$ is immediate. Also, the loss is always at most δ_B , so its expectation has this bound as well. Finally, $\bar{q} \leq q$ and $(m-r)/m \leq 1$. $\square$

Corollary II.153.3 (An explicit anchor budget). *Let $0 < \sigma < 1$ and suppose every row has at most $q \geq 1$ classes. The choice*

$$r = \min \left\{ m, \left\lceil \frac{tq}{\sigma} \right\rceil \right\} \quad (\text{II.153.4})$$

permits $D_C(R) \leq \sigma$. A sufficient budget using the finite-population correction is

$$r \geq \left\lceil \frac{m(tq - \sigma)}{tq + \sigma m} \right\rceil, \quad 0 \leq r \leq m. \quad (\text{II.153.5})$$

The same formulas are valid with q replaced by any upper bound on $\bar{q}$, using a maximum with zero in (II.153.5) if its numerator is negative.

Proof. If (II.153.4) gives $r = m$, the loss is zero. Otherwise $tq/(r+1) < \sigma$. For (II.153.5), rearrange $tq(m-r) \leq \sigma m(r+1)$ to obtain $m(tq - \sigma) \leq r(tq + \sigma m)$, and apply Theorem II.153.2. $\square$

Example II.153.4 (Singletons are not an exception). If $t \geq 2$ and $C = \{x\}$, its class can never contain t distinct anchors. Nevertheless x is certified when $x \in R$, and its exact expected loss is

$$\ell_{m,r,t}(1) = \frac{m-r}{m^2}.$$

No false two-hit estimate is used. If every row consists of at most q singletons, the total expected loss is at most $q(m-r)/m^2$, which can be much smaller than the universal bound.

Remark II.153.5 (Why this improves the power of the error parameter). A truncation proof discards classes of density below σ/q and then requires each surviving class to fail with probability at most σ . That two-stage requirement can introduce an unnecessary second small parameter. Here the loss of a class is multiplied by its own mass before any estimate is made. The exact double count controls the *sum of discarded masses*, directly giving q/r rather than a per-class accuracy demand.

For affine interpolation, take $t = 2$. The repository's corrected long-column lemma requires $6q \leq r\sigma^2$, together with a size condition, and loses at most 2σ ambient mass [175]. Its short-column wrapper removes the separate long-column hypothesis by sampling all columns when needed [176]. Theorem II.153.2 instead uses the one uniform-subset procedure at all column lengths, gives loss σ with (II.153.4), and supplies the same sampled-or-anchored type of output. This comparison concerns the actual stated budgets, not a claim that all possible refinements of the older proof require a quadratic budget.

II.154 The exact finite minimax problem

Source 39, delivered Section 4.

Definition II.154.1 (Worst best-sample loss). For $m \geq 1$, $q \geq 1$, $0 \leq r \leq m$, and $t \geq 1$, let

$$\mathfrak{M}_{m,q,r,t} = \sup_{\mathcal{C}} \min_{\substack{R \subseteq J \\ |R|=r}} D_{\mathcal{C}}(R),$$

where $\mathcal{C}$ ranges over all finite weighted row families with at most q pairwise-disjoint classes in each row, on an m -element column set.

Theorem II.154.2 (Exact finite minimax formula). *The supremum in Definition II.154.1 is attained by a finite unweighted family. Its value is*

$$\boxed{\mathfrak{M}_{m,q,r,t} = \max_{\substack{(s_1, \dots, s_q) \in \mathbb{N}_0^q \\ \sum_i s_i \leq m}} \sum_{i=1}^q \ell_{m,r,t}(s_i).} \quad (\text{II.154.1})$$

An attaining family can be obtained by applying all permutations of J to any maximizing disjoint class profile.

Proof. Let the maximum on the right be A . There are finitely many profiles, so it is attained. Pad rows with empty classes to give exactly q labels. In each row, $\sum_i |C_{h,i}| \leq m$, hence $\sum_i \ell_{m,r,t}(|C_{h,i}|) \leq A$. Averaging over rows and applying (II.153.2) gives an r -set with loss at most A . Thus $\mathfrak{M}_{m,q,r,t} \leq A$.

For the reverse inequality, choose a maximizing profile (s_i) and disjoint sets $C_1, \dots, C_q \subseteq J$ of those sizes. Let the row set be the symmetric group $\text{Sym}(J)$, with equal weights. In row π , use classes $\pi(C_1), \dots, \pi(C_q)$. Fix any r -subset R . The loss in this row is the loss of the original profile against $\pi^{-1}(R)$, divided by m . As π runs uniformly through the symmetric group, $\pi^{-1}(R)$ is uniform among all r -subsets: each such subset occurs $r!(m-r)!$ times. Therefore the average row loss for this fixed R is $\sum_i \ell_{m,r,t}(s_i) = A$.

This holds for every R , so even the best sample has loss A . The family is finite and unweighted, proving both the lower bound and attainment. Equal copies of a row may be removed by passing to the orbit of the labelled partition, since orbit fibres have equal cardinality. $\square$

Corollary II.154.3 (A finite criterion for any proposed universal bound). *A number a bounds the best shared-sample loss for every row family with the parameters (m, q, r, t) if and only if*

$$\sum_{i=1}^q \ell_{m,r,t}(s_i) \leq a \quad \text{for every integer profile with } \sum_i s_i \leq m.$$

Thus any improvement in a universal finite sampling constant can be checked against a finite one-row profile optimization, even though the original minimax problem permits arbitrarily many rows.

II.154.1 Exact evaluation by dynamic programming

For $0 \leq a \leq q$ and $0 \leq b \leq m$, let $V_a(b)$ be the best sum of losses from a classes using at most b columns. Then

$$V_0(b) = 0, \quad V_a(b) = \max_{0 \leq s \leq b} (V_{a-1}(b-s) + \ell_{m,r,t}(s)). \quad (\text{II.154.2})$$

The minimax value is $V_q(m)$. Storing a maximizing s at each state recovers a profile. Once the $m+1$ rational values of ℓ are available, the dynamic program performs $O(qm^2)$ rational comparisons and additions. Memory is $O(qm)$ with parent pointers, or $O(m)$ when only the value is needed. These are arithmetic-operation counts; they do not bound the bit complexity of large binomial coefficients by a constant.

Example II.154.4 (Exact small extremizers). For $(m, q, r, t) = (6, 2, 2, 2)$, the profile $(3, 3)$ gives $\mathfrak{M} = 3/5$. The 20 labelled partitions of a six-element set into two three-element classes form a family for which all 15 two-element anchor sets have loss exactly $3/5$.

For $(7, 2, 3, 2)$, the maximizing profile $(3, 3)$ leaves one unused column and gives $\mathfrak{M} = 96/245$. Its 140 labelled placements have this average loss against every one of the 35 three-element samples. These identities are also checked by exhaustive rational enumeration in the supplied certificates.

Proposition II.154.5 (Boundary regimes). *For $q \geq 1$, one has $\mathfrak{M}_{m,q,m,t} = 0$. If $r < t$, then*

$$\mathfrak{M}_{m,q,r,t} = 1 - \frac{r}{m}. \quad (\text{II.154.3})$$

Proof. The first assertion follows because every point is sampled. If $r < t$, no class can have t sampled points, so only sampled targets are retained. For a class of size s , its expected discarded count is $s(1 - r/m)$; hence $\ell_{m,r,t}(s) = s(m - r)/m^2$. The profile maximum is obtained by using all m columns, for example in a single class. This gives (II.154.3). $\square$

Remark II.154.6 (The scope of the extremal family). The symmetrization theorem is valid with no field structure on J . If a field structure is subsequently added, the class family still need not exhibit the horizontal regularity required elsewhere in Gowers's proof. Moreover, even if a function is assigned to each class, a point left uncertified by this rule might be algebraically recoverable in some other way. The theorem proves minimax optimality of the rule in Definition II.152.2, and no stronger algorithmic lower bound.

II.155 A closed finite profile formula and the affine quadratic cap

Source 39, delivered Section 5.

The profile optimization in Theorem II.154.2 can in fact be solved without dynamic programming. The key extra fact is that the one-class loss is concave up to its first maximum and never increases afterwards. The hypergeometric expression makes this shape accessible through exact finite differences.

Write $\Delta u(s) = u(s+1) - u(s)$ and $\Delta^2 u(s) = u(s+2) - 2u(s+1) + u(s)$. These expressions are only used when their arguments lie between 0 and m .

Lemma II.155.1 (Finite differences of a hypergeometric average). *Let g be any function on $\{0, \dots, n\}$ and let $A_n(s) = \mathbb{E}[g(Z_s)]$ for $Z_s \sim \text{Hyp}(m, s, n)$, with $0 \leq n \leq m$. For $n \geq 1$,*

$$\Delta A_n(s) = \frac{n}{m} \mathbb{E}[\Delta g(Y)], \quad Y \sim \text{Hyp}(m-1, s, n-1). \quad (\text{II.155.1})$$

For $n \geq 2$ and $m \geq 2$,

$$\Delta^2 A_n(s) = \frac{n(n-1)}{m(m-1)} \mathbb{E}[\Delta^2 g(Y)], \quad Y \sim \text{Hyp}(m-2, s, n-2). \quad (\text{II.155.2})$$

Proof. Couple the class of size s with that of size $s+1$ by adding one distinguished column. A uniform n -subset contains that column with probability n/m . Conditional on containing it, its other $n-1$ columns are uniform in the remaining population of size $m-1$, with s successes. If the distinguished column is absent, the class count does not change; if it is present, the change in g is $\Delta g(Y)$. This proves (II.155.1). Apply the same argument to the hypergeometric average on its right, with population $m-1$, sample size $n-1$, and function Δg , to obtain (II.155.2). $\square$

Lemma II.155.2 (Exact curvature and the one-class shape). *Fix m, r, t . The sequence $s \mapsto \ell_{m,r,t}(s)$ is nondecreasing up to its first maximizing index $a = a_{m,r,t}$ and nonincreasing thereafter. It is discretely concave on $\{0, \dots, a\}$.*

More explicitly, when $1 \leq t \leq r < m$ and $0 \leq s \leq m-2$, put

$$K_{m,r} = \frac{r(m-r)}{m^2(m-1)\binom{m-2}{r-1}} > 0.$$

Then

$$\begin{aligned} \Delta^2 \ell_{m,r,t}(s) = K_{m,r} & \left[t \binom{s}{t} \binom{m-s-2}{r-t-1} \right. \\ & \left. - (t+1) \binom{s}{t-1} \binom{m-s-2}{r-t} \right]. \end{aligned} \quad (\text{II.155.3})$$

For $r > t$ this factors as

$$\Delta^2 \ell_{m,r,t}(s) = \frac{K_{m,r}}{r-t} \binom{s}{t-1} \binom{m-s-2}{r-t-1} [(r+1)(s+2) - (t+1)(m+1)]. \quad (\text{II.155.4})$$

For $r = t$ the second difference is nonpositive everywhere.

Proof. If $r = m$, the sequence is identically zero. If $r < t$, then $\ell(s) = s(m-r)/m^2$ by Proposition II.154.5, so it is linear and nondecreasing. These cases already give the claimed shape.

Suppose $1 \leq t \leq r < m$. In (II.153.1) set $g(z) = z \mathbf{1}_{\{z \leq t\}}$ for $z \geq 0$. Its second difference is zero except at $z = t-1$ and $z = t$, where it equals $-(t+1)$ and t , respectively. Apply Lemma II.155.1 with

$n = r + 1$, and multiply by the prefactor $(m - r)/(m(r + 1))$. The remaining hypergeometric variable has population $m - 2$, class size s , and sample size $r - 1$. Writing its two relevant probabilities with binomial coefficients yields (II.155.3).

If $r = t$, its first term is zero, proving global concavity. If $r > t$, use

$$t \binom{s}{t} = (s - t + 1) \binom{s}{t - 1}, \quad \binom{m - s - 2}{r - t} = \frac{m - s - r + t - 1}{r - t} \binom{m - s - 2}{r - t - 1}.$$

These identities remain valid at the boundary with the zero convention. The bracket after factoring is

$$(s - t + 1)(r - t) - (t + 1)(m - s - r + t - 1) = (r + 1)(s + 2) - (t + 1)(m + 1),$$

which proves (II.155.4).

Consequently the finite-difference sequence $d_s = \ell(s + 1) - \ell(s)$ is first nonincreasing and then nondecreasing; in the case $r = t$ it is simply nonincreasing. Also $d_0 = (m - r)/m^2 > 0$, while $\ell(m) = 0$ and $\ell(m - 1) \geq 0$, so $d_{m-1} \leq 0$. Every term on the final nondecreasing part is therefore nonpositive. On the initial nonincreasing part, the differences can cross from positive to nonpositive at most once. It follows that ℓ increases to its first maximum and never increases after it. The differences before that first maximum belong to the nonincreasing part, proving discrete concavity up to a . $\square$

Theorem II.155.3 (Balanced-and-capped finite extremizers). *Let $a = a_{m,r,t}$ be the least size maximizing $\ell_{m,r,t}$. Define*

$$B = \min\{m, qa\}, \quad v = \lfloor B/q \rfloor, \quad u = \lceil B/q \rceil, \quad \beta = B - qv. \quad (\text{II.155.5})$$

Then

$$\boxed{\mathfrak{M}_{m,q,r,t} = (q - \beta)\ell_{m,r,t}(v) + \beta\ell_{m,r,t}(u).} \quad (\text{II.155.6})$$

An extremizing profile has $q - \beta$ classes of size v and β classes of size u . In particular, an extremizer exists with at most two positive class sizes, and those sizes differ by at most one.

Proof. Start with any feasible profile. Replace every size greater than a by a . By Lemma II.155.2, this does not decrease its objective and does not increase its total size. All sizes now lie in $[0, a]$. Since ℓ is nondecreasing there, add columns to classes until their total is $B = \min\{m, qa\}$; this is feasible and does not decrease the objective.

If two sizes satisfy $y \geq x + 2$, move one column from the larger to the smaller. The objective changes by

$$\ell(x + 1) - \ell(x) - (\ell(y) - \ell(y - 1)) = \Delta\ell(x) - \Delta\ell(y - 1) \geq 0,$$

using discrete concavity on $[0, a]$. Iterating this balancing operation terminates, for example because the sum of squared sizes strictly decreases at each nontrivial step. At termination the sizes differ by at most one and sum to B , hence have the multiplicities in (II.155.5). Their objective is (II.155.6). This profile is feasible, so Theorem II.154.2 proves equality. $\square$

[write] **6 October 2026; sources 34–39.** *Source 19's question, answered for this rule only* (checked at the write by exact rational arithmetic; the intake found the same counts). Source 19's Question 1 asks for the structure of the maximizers of its finite minimax (item II.39). For source 39's rule the formula (II.155.6) agrees with an exact maximization over all class profiles in all 2,835 cases $m \leq 18$, $0 \leq r \leq m$, $t \leq 3$, $q \leq 5$. For source 19's rule it does not: with source 19's one-class loss $W_{m,r,t}(s) = (s/m) \Pr\{X < \min(t, s)\}$ in place

of $\ell_{m,r,t}$, the balanced-and-capped value differs from the exact maximum in 199 of the same 2,835 cases; the first is $m = 2, r = 1, t = 2, q = 2$, where the formula gives $1/2$ and the exact value is 1 (one class of size 2, lost with its one sampled point; under source 39's rule the same profile loses $1/2$). Source 19's one-class loss is not concave up to its first maximum, which is what the proof of Theorem II.155.3 uses. The rules also give different values, for example 0.3512 (source 39) against 0.3944 (source 19) at $m = 20, q = 2, r = 4, t = 2$. So the classification for source 19's rule stays open (item II.39).

When $\beta = 0$, one has $u = v$, so formula (II.155.6) never requires evaluating ℓ outside its domain. When $r = m$, the first mode is $a = 0$, giving zero. When $r < t$, it is $a = m$, and the formula reproduces $1 - r/m$. Empty classes cause no problem when $q > m$.

This theorem solves the finite optimization, rather than merely reducing it to a dynamic program. A scan of the $m + 1$ one-class values determines a ; the remaining computation is a single balancing operation. The dynamic program in Section II.154 is retained as an independent exact verification method.

Theorem II.155.4 (An explicit quadratic cap for affine fibres). *Let $t = 2$. If $r = m$, then $a_{m,r,2} = 0$; if $r < 2$ and $r < m$, then $a_{m,r,2} = m$. For $2 \leq r < m$, define*

$$\zeta_{m,r} = \frac{B_{m,r} + \sqrt{B_{m,r}^2 + 4(r^2 - 1)(m - r)(m - r + 1)}}{2(r^2 - 1)}, \quad B_{m,r} = (r - 2)m + 2r - 1. \quad (\text{II.155.7})$$

Then the least maximizing class size is exactly

$$a_{m,r,2} = \lceil \zeta_{m,r} \rceil. \quad (\text{II.155.8})$$

Together with Theorem II.155.3, this is a closed finite formula for the affine minimax loss, requiring no search over profiles or class sizes.

Proof. The endpoint and $r < 2$ cases have already been established. For $2 \leq r < m$, apply (II.155.1) to (II.153.1). Since $\Delta g(z) = 1$ for $z = 0, 1$, $\Delta g(2) = -2$, and it is zero afterwards,

$$\Delta \ell_{m,r,2}(s) = \frac{m - r}{m^2} [\mathbb{P}(Z = 0) + \mathbb{P}(Z = 1) - 2\mathbb{P}(Z = 2)], \quad Z \sim \text{Hyp}(m - 1, s, r).$$

Apart from positive factors, the bracket has numerator

$$\binom{m - 1 - s}{r} + s \binom{m - 1 - s}{r - 1} - s(s - 1) \binom{m - 1 - s}{r - 2}.$$

For $0 \leq s \leq m - r + 1$, factor out the positive quantity $\binom{m - 1 - s}{r - 2} / (r(r - 1))$. The remaining quadratic is

$$-(r^2 - 1)s^2 + B_{m,r}s + (m - r)(m - r + 1). \quad (\text{II.155.9})$$

Its constant term is positive and its quadratic coefficient negative, so it has one positive root, namely (II.155.7). At $s = m - r + 1 \geq 2$, the unfactored numerator is strictly negative: only its last term is nonzero. Thus the positive root lies strictly between 0 and $m - r + 1$.

The first difference is positive at integer $s < \zeta_{m,r}$ and nonpositive at the first integer $s \geq \zeta_{m,r}$. By the unimodality already proved, this first change determines the least maximum, giving $a_{m,r,2} = \lceil \zeta_{m,r} \rceil$. If the root is an integer, the difference there is zero and the first of the two equal maximizing values is precisely that integer, so the ceiling formula still applies. $\square$

Remark II.155.5 (Exact implementation of the radical ceiling). Let $A = r^2 - 1$, $B = B_{m,r}$, $C = (m - r)(m - r + 1)$, and $D = B^2 + 4AC$. Set $k = \lfloor (B + \lfloor \sqrt{D} \rfloor) / (2A) \rfloor$. This is $\lfloor \zeta_{m,r} \rfloor$. The cap equals k when $Ak^2 - Bk - C = 0$, and otherwise equals $k + 1$. Integer square root and this exact equality test avoid floating-point decisions near an integer root. The supplied implementation uses this procedure.

Example II.155.6 (A genuinely finite improvement). For $(m, q, r, t) = (20, 2, 4, 2)$, the affine root formula gives $a = 7$. Since $qa = 14 \leq 20$, the exact worst-case profile is $(7, 7)$ and

$$\mathfrak{M}_{20,2,4,2} = \frac{1001}{2850} = 0.351228\dots$$

The simple universal bound is $16/25 = 0.64$. The improved value is not an asymptotic estimate: every such row family has a four-anchor sample with loss at most $1001/2850$, and a finite permutation-orbit family attains equality.

II.156 The Poisson profile and its exact optimization

Source 39, delivered Section 6.

The expression F_t from (F.44.4) has a direct interpretation:

$$F_t(x) = x\mathbb{P}\{\text{Pois}(x) < t\}. \quad (\text{II.156.1})$$

It balances class mass against the probability that a class receives too few anchors. This section solves the associated continuous optimization exactly; Section II.157 justifies its use for the finite minimax problem.

Lemma II.156.1 (Unique maximum and the useful concavity interval). *For every $t \geq 1$, the function F_t has a unique positive maximizer λ_t . It is characterized by*

$$\sum_{j=0}^{t-1} \frac{\lambda_t^j}{j!} = \frac{\lambda_t^t}{(t-1)!}. \quad (\text{II.156.2})$$

One has $0 < \lambda_t \leq t$, with strict inequality when $t > 1$. The function is strictly increasing before λ_t , strictly decreasing after λ_t , and concave on $[0, \lambda_t]$.

Proof. Put $P_{t-1}(x) = \sum_{j=0}^{t-1} x^j / j!$. Direct differentiation gives

$$F'_t(x) = e^{-x} \left(P_{t-1}(x) - \frac{x^t}{(t-1)!} \right), \quad (\text{II.156.3})$$

$$F''_t(x) = \frac{e^{-x} x^{t-1}}{(t-1)!} (x - t - 1). \quad (\text{II.156.4})$$

For $x > 0$, the equation $F'_t(x) = 0$ becomes

$$\sum_{j=0}^{t-1} \frac{x^{j-t}}{j!} = \frac{1}{(t-1)!}.$$

The left side is strictly decreasing from infinity to zero. This proves a unique root and the claimed signs of the derivative.

At $x = t$, the terms $t^j/j!$ are nondecreasing for $0 \leq j \leq t-1$, and each is at most $t^{t-1}/(t-1)!$. Therefore $P_{t-1}(t) \leq t^t/(t-1)!$, strictly for $t > 1$. Hence $\lambda_t \leq t$, strictly for $t > 1$. Formula (II.156.4) now proves concavity on $[0, \lambda_t]$. $\square$

Theorem II.156.2 (Poissonized subprobability optimization). *For real $r > 0$, integer $q \geq 1$, and integer $t \geq 1$,*

$$\max_{\substack{p_1, \dots, p_q \geq 0 \\ \sum_i p_i \leq 1}} \sum_{i=1}^q p_i \mathbb{P}\{\text{Pois}(rp_i) < t\} = \frac{q}{r} F_t \left(\min \left\{ \lambda_t, \frac{r}{q} \right\} \right). \quad (\text{II.156.5})$$

An optimizer has all rp_i equal to $\min\{\lambda_t, r/q\}$.

Proof. Write $x_i = rp_i$. The optimization is $r^{-1} \max \sum_i F_t(x_i)$ subject to $x_i \geq 0$ and $\sum_i x_i \leq r$. Any coordinate larger than λ_t may be replaced by λ_t , decreasing the resource consumption and increasing the objective. We may therefore restrict to $x_i \in [0, \lambda_t]$.

By concavity and Jensen's inequality,

$$\sum_{i=1}^q F_t(x_i) \leq q F_t \left(\frac{1}{q} \sum_i x_i \right).$$

The average is at most both λ_t and r/q . Since F_t is increasing on $[0, \lambda_t]$, this is at most $q F_t(\min\{\lambda_t, r/q\})$. Equal coordinates at that value are feasible and attain the bound. $\square$

The inequality $\sum p_i \leq 1$, rather than equality, is part of the theorem. If $r/q > \lambda_t$, the optimizer leaves mass $1 - q\lambda_t/r$ unused. If the classes were forced to partition all columns with no leftover domain, the optimization would be different. The partial-graph application permits the subprobability formulation naturally.

Corollary II.156.3 (The affine constant). *For $t = 2$,*

$$\lambda_2 = \varphi, \quad \kappa_2 = \varphi^3 e^{-\varphi} = 0.839962094657 \dots$$

For $t = 1$, $\lambda_1 = 1$ and $\kappa_1 = e^{-1}$.

Proof. Equation (II.156.2) is $1 + x = x^2$ for $t = 2$. Its positive root is φ , and $F_2(\varphi) = \varphi(1 + \varphi)e^{-\varphi} = \varphi^3 e^{-\varphi}$. The $t = 1$ case is immediate. $\square$

[write] **6 October 2026; sources 34–39.** This is source 19's Corollary II.69.5 (the constant 0.839962094657; source 19's Lemma II.69.1 gives the maximizing mean). In the regime $q/r \rightarrow 0$ the two rules have the same limit, since the sampled points of a failed class carry vanishing mass.

II.157 Uniform asymptotics and a sharp sample-complexity law

Source 39, delivered Section 7.

Pointwise Poisson convergence is not by itself enough to optimize over every class-size profile. The maximizing sizes vary with m , r , and q , and a profile may contain very large classes. We therefore prove a uniform marginal estimate, including a tail bound, before interchanging the limit with the profile maximum.

Table II.156.1: Poisson thresholds and sharp sparse-sampling coefficients. Decimals are numerical approximations; (II.156.2) defines the exact values.

t	λ_t	κ_t	transition $1/\lambda_t$
1	1.0000000000	0.3678794412	1.0000000000
2	1.6180339887	0.8399620947	0.6180339887
3	2.2695308421	1.3711016049	0.4406197005
4	2.9451861612	1.9423809381	0.3395371108
5	3.6395471265	2.5435343541	0.2747594591
6	4.3490476051	3.1681848157	0.2299353998
7	5.0711843460	3.8120212302	0.1971925948
8	5.8041103933	4.4719539621	0.1722916920

Lemma II.157.1 (An elementary product bound for sampling without replacement). *Let $z_1, \dots, z_m \geq 0$ and let R be a uniform r -subset of $\{1, \dots, m\}$. Then*

$$\mathbb{E} \prod_{i \in R} z_i \leq \left(\frac{1}{m} \sum_{i=1}^m z_i \right)^r. \quad (\text{II.157.1})$$

In particular, for $A \sim \text{Hyp}(m, s, r)$ and $c > 0$,

$$\mathbb{E} e^{-cA} \leq \left(1 - \frac{s}{m} (1 - e^{-c}) \right)^r \leq \exp \left(-\frac{rs}{m} (1 - e^{-c}) \right). \quad (\text{II.157.2})$$

Proof. The left side of (II.157.1) is the normalized elementary symmetric polynomial $e_r(z_1, \dots, z_m) / \binom{m}{r}$. Holding all but two coordinates u, v fixed, this polynomial before normalization has the form $A_0 + (u+v)B_0 + uvC_0$ with $A_0, B_0, C_0 \geq 0$. Replacing u, v by their average preserves their sum and increases their product, so does not decrease e_r .

Repeatedly average a largest and a smallest coordinate. The sum remains fixed and the sum of squared deviations from the mean decreases by half the square of the current spread. Thus the spread tends to zero, and the coordinates converge to the common mean. By continuity, e_r at the original vector is at most $\binom{m}{r}$ times the r th power of the mean. The cases $r = 0, 1$ also follow directly.

For (II.157.2), set $z_i = e^{-c}$ on the class and $z_i = 1$ outside it. Apply (II.157.1) and $1 - u \leq e^{-u}$. $\square$

Theorem II.157.2 (Uniform one-class Poisson limit). *Fix $t \geq 1$. Along any sequence of integer pairs (m, r) with $r \rightarrow \infty$ and $r/m \rightarrow 0$,*

$$\sup_{0 \leq s \leq m} \left| r \ell_{m,r,t}(s) - F_t \left(\frac{rs}{m} \right) \right| \rightarrow 0. \quad (\text{II.157.3})$$

The supremum is over integer class sizes.

Proof. Write $x = rs/m$ and $A \sim \text{Hyp}(m, s, r)$. By the definition of ℓ ,

$$r \ell_{m,r,t}(s) = x \sum_{j=0}^{t-1} \left(1 - \frac{j}{s} \right) \mathbb{P}(A = j) \quad (\text{II.157.4})$$

when $s > 0$, with zero terms for impossible j . The $s = 0$ case is zero. The expression is nonnegative and at most $x\mathbb{P}(A < t)$.

First consider any sequence with $x \rightarrow x_0 \in (0, \infty)$. Then $s \rightarrow \infty$, $s/m \rightarrow 0$, and eventually $s + r < m$. The zero-count probability satisfies

$$\mathbb{P}(A = 0) = \prod_{i=0}^{r-1} \left(1 - \frac{s}{m-i}\right),$$

$$\log \mathbb{P}(A = 0) = -x + O\left(\frac{xr}{m} + \frac{x^2}{r}\right).$$

Indeed, uniformly for $0 \leq i < r$, the quantity $s/(m-i)$ is $O(x/r)$; summing its first-order terms gives $x + O(xr/m)$, and summing its squares gives $O(x^2/r)$. For any fixed $j \geq 0$,

$$\frac{\mathbb{P}(A = j)}{\mathbb{P}(A = 0)} = \frac{(s)_j (r)_j}{j! \prod_{a=1}^j (m - s - r + a)} \rightarrow \frac{x_0^j}{j!}, \quad (\text{II.157.5})$$

where $(s)_j = s(s-1)\cdots(s-j+1)$. Consequently $\mathbb{P}(A = j) \rightarrow e^{-x_0} x_0^j / j!$, and (II.157.4) tends to $F_t(x_0)$. If $x \rightarrow 0$, both $r\ell$ and $F_t(x)$ lie between zero and x , so the same conclusion holds.

These sequential statements imply uniform convergence on every bounded interval of x . Otherwise one could choose a sequence of bounded x at which the error stays above a fixed positive number, pass to a convergent subsequence, and contradict the preceding argument.

It remains to control unbounded x . Since $A < t$ means $A \leq t-1$,

$$\mathbb{P}(A < t) \leq e^{t-1} \mathbb{E}e^{-A} \leq \exp(t-1 - (1-e^{-1})x)$$

by Lemma II.157.1. Thus, uniformly in m, r, s ,

$$0 \leq r\ell_{m,r,t}(s) \leq x \exp(t-1 - (1-e^{-1})x). \quad (\text{II.157.6})$$

The same bound holds for $F_t(x)$ by the Laplace transform of a Poisson variable, or it follows directly from (II.156.1). Both tails tend uniformly to zero as $x \rightarrow \infty$. Combining the compact interval convergence with these tail estimates proves (II.157.3). $\square$

Theorem II.157.3 (Asymptotic minimax profile). *Fix $t \geq 1$. Let m, r, q vary with $r \rightarrow \infty$, $r/m \rightarrow 0$, and $q \geq 1$. If q/r stays bounded, then*

$$\mathfrak{M}_{m,q,r,t} = \frac{q}{r} \left[F_t\left(\min\left\{\lambda_t, \frac{r}{q}\right\}\right) + o(1) \right]. \quad (\text{II.157.7})$$

In particular, if $q/r \rightarrow 0$, then $\mathfrak{M}_{m,q,r,t} \sim \kappa_t q/r$.

Proof. Let $\varepsilon_{m,r}$ denote the supremum in (II.157.3). It tends to zero. For any feasible finite profile (s_i) , put $x_i = rs_i/m$. Then $\sum_i x_i \leq r$, so Theorems II.154.2, II.156.2, and II.157.2 give

$$\mathfrak{M}_{m,q,r,t} \leq \frac{q}{r} \left[F_t\left(\min\left\{\lambda_t, \frac{r}{q}\right\}\right) + \varepsilon_{m,r} \right].$$

For a matching lower bound, set $x_* = \min\{\lambda_t, r/q\}$ and choose all q class sizes equal to $s_* = \lfloor mx_*/r \rfloor$. This is feasible since $qs_* \leq m x_* \leq m$. Also $0 \leq x_* - rs_*/m < r/m \rightarrow 0$. The function F_t is uniformly continuous on $[0, \lambda_t]$, so

$$\mathfrak{M}_{m,q,r,t} \geq q\ell_{m,r,t}(s_*) \geq \frac{q}{r} (F_t(x_*) - o(1) - \varepsilon_{m,r}).$$

This proves (II.157.7). If $q/r \rightarrow 0$, then $x_* = \lambda_t$ eventually, and $\kappa_t > 0$, giving the relative asymptotic. $\square$

[write] **6 October 2026; sources 34–39.** Source 19's Theorem II.69.4 is the case $q/r \rightarrow 0$; the regime q/r bounded, where it is best to use only part of the column set, is new here and advances source 19's question on rates in the sharp law (item II.39).

Corollary II.157.4 (Two asymptotic regimes). *Under the hypotheses above, if $q/r \rightarrow \theta \in (0, \infty)$, then*

$$\mathfrak{M}_{m,q,r,t} \longrightarrow \Psi_t(\theta) := \begin{cases} \theta \kappa_t, & 0 < \theta \leq 1/\lambda_t, \\ e^{-1/\theta} \sum_{j=0}^{t-1} \frac{\theta^{-j}}{j!}, & \theta \geq 1/\lambda_t. \end{cases} \quad (\text{II.157.8})$$

In the first regime, asymptotically extremal class sizes are $s_i \sim \lambda_t m/r$ and need not exhaust the columns. In the second, they are $s_i \sim m/q$ and asymptotically exhaust the column budget.

Proof. Substitute $r/q \rightarrow 1/\theta$ in (II.157.7). The profiles in the lower bound proof have the stated sizes. At the transition the two expressions coincide because $\kappa_t = F_t(\lambda_t)$. $\square$

Proposition II.157.5 (Shape of the limiting loss curve). *Extend Ψ_t by $\Psi_t(0) = 0$. It is continuously differentiable, increasing, and concave on $[0, \infty)$. It is linear on $[0, 1/\lambda_t]$ and strictly concave thereafter. Moreover, $\Psi_t(\theta) \rightarrow 1$ as $\theta \rightarrow \infty$.*

Proof. On the nonlinear branch, differentiation gives

$$\Psi'_t(\theta) = \frac{e^{-1/\theta}}{(t-1)! \theta^{t+1}} > 0, \quad \frac{\Psi''_t(\theta)}{\Psi'_t(\theta)} = \frac{1 - (t+1)\theta}{\theta^2} < 0.$$

The last inequality holds because $\theta \geq 1/\lambda_t \geq 1/t$. At $\theta = 1/\lambda_t$, the right derivative is $e^{-\lambda_t} \lambda_t^{t+1}/(t-1)! = \kappa_t$ by (II.156.2); this equals the slope of the linear branch. The limit at infinity follows from (II.157.8). $\square$

The statement about $\theta \rightarrow \infty$ describes the explicitly defined curve Ψ_t , not an additional uniform finite-parameter limit outside the stated hypotheses of Theorem II.157.3.

Theorem II.157.6 (Sharp sparse-sampling anchor budget). *For $0 < \sigma < 1$, define the smallest worst-case sufficient sample size by*

$$r_*(m, q, t, \sigma) = \min\{0 \leq r \leq m : \mathfrak{M}_{m,q,r,t} \leq \sigma\}.$$

Fix t . Along sequences with $\sigma \rightarrow 0$, $q \geq 1$, and $q/(\sigma m) \rightarrow 0$, one has

$$r_*(m, q, t, \sigma) \sim \frac{\kappa_t q}{\sigma}. \quad (\text{II.157.9})$$

For affine certification the coefficient is $\kappa_2 = \varphi^3 e^{-\varphi}$.

Proof. The minimum exists since sampling all m columns gives zero loss. Fix $0 < \epsilon < \kappa_t$ and compare

$$r_- = \left\lfloor \frac{(\kappa_t - \epsilon)q}{\sigma} \right\rfloor, \quad r_+ = \left\lceil \frac{(\kappa_t + \epsilon)q}{\sigma} \right\rceil.$$

Both tend to infinity, both satisfy $r_{\pm}/m \rightarrow 0$, and $q/r_{\pm} \rightarrow 0$. Theorem II.157.3 therefore gives

$$\frac{\mathfrak{M}_{m,q,r_-,t}}{\sigma} \rightarrow \frac{\kappa_t}{\kappa_t - \epsilon} > 1, \quad \frac{\mathfrak{M}_{m,q,r_+,t}}{\sigma} \rightarrow \frac{\kappa_t}{\kappa_t + \epsilon} < 1.$$

The best-sample loss is nonincreasing in r , by Remark II.152.3, and the same is true of its supremum. Hence eventually $r_- < r_* \leq r_+$. Divide by q/σ and let $\epsilon \downarrow 0$. $\square$

Equation (II.157.9) is a genuine optimality statement for a sample-complexity problem. It does not say that the simple finite coefficient t in Corollary II.153.3 is sharp, nor that the asymptotic coefficient κ_t can be substituted into that finite corollary without an error term and additional scale assumptions.

II.158 Deterministic anchor selection by conditional expectation

Source 39, delivered Section 8.

The existence proof can be made constructive when the row classes are given explicitly. There is no need to enumerate all $\binom{m}{r}$ possible final samples.

Suppose $T \subseteq J$ is a partial sample of size $a \leq r$. Complete it by a uniform $(r-a)$ -subset of $J \setminus T$. For a class C , write

$$b = |C \cap T|, \quad c = |C \setminus T|, \quad n = m - a, \quad u = r - a.$$

If $b \geq t$, its conditional expected discarded count is zero. Otherwise it is

$$L_C(T) = \frac{1}{\binom{n}{u}} \sum_{j=0}^{t-b-1} (c-j) \binom{c}{j} \binom{n-c}{u-j}, \quad (\text{II.158.1})$$

with zero terms outside the feasible range. Indeed, the additional hits have law $\text{Hyp}(n, c, u)$, and after j such hits there are $c-j$ unsampled points in the class.

Define the conditional potential

$$\Phi(T) = \frac{1}{mW} \sum_h w_h \sum_i L_{C_{h,i}}(T). \quad (\text{II.158.2})$$

It is the expected final loss conditioned on T being contained in the final sample, with the remaining columns chosen uniformly.

Lemma II.158.1 (One-step potential identity). *If $|T| < r$, then*

$$\Phi(T) = \frac{1}{m - |T|} \sum_{x \in J \setminus T} \Phi(T \cup \{x\}). \quad (\text{II.158.3})$$

Proof. Choose the next column uniformly from $J \setminus T$, then choose the rest uniformly without replacement. This yields a uniform final r -subset containing T . More explicitly, each particular completion of size $r - |T|$ can have any of its $r - |T|$ points selected first, and these choices give the same total probability for every completion. The law of total expectation now gives (II.158.3). $\square$

Theorem II.158.2 (Exact derandomization). *Start with $T = \emptyset$. While $|T| < r$, add a column minimizing $\Phi(T \cup \{x\})$ over $x \notin T$. The final sample R satisfies*

$$D_C(R) = \Phi(R) \leq \Phi(\emptyset) = \frac{1}{W} \sum_h w_h \sum_i \ell_{m,r,t}(|C_{h,i}|). \quad (\text{II.158.4})$$

For rational row weights, every comparison can be performed exactly with rational arithmetic.

Proof. A minimum of finitely many real numbers is at most their average, so Lemma II.158.1 makes the potential nonincreasing at every step. When $|T| = r$, there is no remaining randomness; formula (II.158.1) is the actual discarded count. At the initial state it is the unconditional expectation from Lemma II.153.1. $\square$

The supplied implementation follows this algorithm with exact rational arithmetic. Its main module is `code/anchor_selection.py`; values have Python's `fractions.Fraction` type. It validates that classes within a row are disjoint and that columns and parameters are in range. A straightforward implementation evaluates at most rm candidate extensions, each using the supplied row classes. This is a finite algorithm for explicitly represented data; it is not asserted to be an efficient implementation of the entire inverse proof, where the class descriptions and the number of rows may be extremely large.

For a fixed row family the derandomized loss can be much smaller than the universal minimax value. The permutation-orbit construction is the limiting case: the final loss is the same for every sample, so no deterministic choice can improve on the average. This gives a useful exact test case for the implementation.

II.159 Polynomial reconstruction, general interpolation spaces, and errors

Source 39, delivered Section 9.

II.159.1 A polynomial graph-cover theorem

Let $\mathbb{F}$ be any field, and now regard J as an m -element subset of $\mathbb{F}$. The distinctness of the elements of J is the only hypothesis needed for the Vandermonde denominators. Let $B \subseteq H \times J$ and let $f : B \rightarrow \mathbb{F}$. Suppose each fibre B_h has a partition into at most q classes such that on each class $f(h, \cdot)$ agrees with a polynomial of degree at most $t - 1$. The polynomials may depend on both the row and the class.

For $a \in J$, define the partial horizontal section

$$B_a = \{h : (h, a) \in B\}, \quad \psi_a(h) = f(h, a) \quad (h \in B_a).$$

For a t -subset $S \subseteq J$ and $a \in S$, put

$$L_{S,a}(x) = \prod_{b \in S \setminus \{a\}} \frac{x - b}{a - b}. \quad (\text{II.159.1})$$

These polynomials satisfy $L_{S,a}(b) = \mathbf{1}_{\{a=b\}}$ for $b \in S$.

Theorem II.159.1 (Shared polynomial interpolants). *There is an r -subset $R \subseteq J$ and an ambient good set E_R with weighted complement at most*

$$\frac{1}{W} \sum_h w_h \sum_i \ell_{m,r,t}(|C_{h,i}|) \leq \frac{tq(m-r)}{m(r+1)}$$

(as a fraction of mW), such that every $(h, x) \in B \cap E_R$ lies on one of the following partial graphs:

$$(h, x) \mapsto \psi_a(h) \quad (a \in R), \quad \text{defined when } h \in B_a; \quad (\text{II.159.2})$$

$$I_S(h, x) = \sum_{a \in S} L_{S,a}(x) \psi_a(h) \quad \left(S \in \binom{R}{t} \right), \quad \text{defined when } h \in \bigcap_{a \in S} B_a. \quad (\text{II.159.3})$$

There are at most $r + \binom{r}{t}$ candidate expressions. Every anchor actually used to certify a nonsampled target belongs to $B \cap E_R$ in the same row.

Proof. Choose the sample by Theorem II.153.2 and use the good set from Definition II.152.2. If $x \in R$, then the single expression with $a = x$ equals $f(h, x)$.

Otherwise, certification means that the class containing x contains at least t elements of R . Choose any t of them and call this set S . Let P be the polynomial of degree at most $t - 1$ agreeing with $f(h, \cdot)$ on the class. Lagrange interpolation gives

$$P(x) = \sum_{a \in S} L_{S,a}(x) P(a) = \sum_{a \in S} L_{S,a}(x) \psi_a(h) = I_S(h, x).$$

For completeness, the interpolation formula follows because its right side is a polynomial of degree at most $t - 1$ agreeing with P at the t distinct points of S ; their difference has t roots and degree at most $t - 1$, and hence is zero. The root bound follows by repeatedly factoring $X - a$ at each root over a field.

The domain conditions hold because S is contained in the same class. All its points are sampled, and the definition of E_R retains every sampled domain point. The number of expressions is the number of singletons plus the number of t -subsets of R . $\square$

For $t = 2$, the paired expression can be written without separating slope and intercept:

$$I_{\{a,b\}}(h, x) = \frac{x-b}{a-b} \psi_a(h) + \frac{a-x}{a-b} \psi_b(h). \quad (\text{II.159.4})$$

This is a linear combination of only two horizontal sections, with affine coefficients in x . If each section is multilinear in the base coordinates, the displayed function is multilinear in the base coordinates and x . There is no need to first place all pairwise slopes and intercepts into a larger common union and then refine that union.

Remark II.159.2 (Partial graphs and harmless spurious candidates). The theorem is a cover theorem, not a uniqueness theorem. Different candidate expressions may coincide, and many candidates need not agree with f at a given point. Only one correct expression is required for each retained point. Later list-cover bounds count all candidates because removing spurious ones is not needed for validity.

II.159.2 Unisolvent reconstruction spaces

Definition II.159.3 (A unisolvent space). A t -dimensional linear subspace $\mathcal{W} \subseteq \mathbb{F}^J$ is called t -point unisolvent if, for every $S \in \binom{J}{t}$, the evaluation map $\mathcal{W} \rightarrow \mathbb{F}^S$ is an isomorphism. We additionally require that the constant function 1 belong to $\mathcal{W}$.

The assumption about constants ensures that the direct single-anchor expressions in (II.159.2) remain in the same reconstruction space. If only the nonsampled targets were being considered, this extra assumption would not be needed.

Corollary II.159.4 (Interpolation-space extension). *Replace polynomial class functions in Theorem II.159.1 by members of a t -point unisolvent space $\mathcal{W}$ containing constants. Then the same sample, mass bounds, and candidate count hold, with the Lagrange functions $L_{S,a}$ replaced by the unique elements of $\mathcal{W}$ whose evaluations on S are the standard coordinate vectors.*

Proof. Invert the evaluation isomorphism on each S . If $P \in \mathcal{W}$, its vector of values on S determines it, so $P = \sum_{a \in S} P(a)L_{S,a}$ in $\mathcal{W}$. The rest of the proof is exactly the certification and domain argument of Theorem II.159.1. Constant functions handle sampled targets. $\square$

The polynomial theorem itself remains valid when $m < t$; in that case no t -subset of J exists, and only direct sampled targets can be certified. A t -dimensional subspace of $\mathbb{F}^J$ cannot exist when $m < t$, so the general space formulation is only applicable when its hypotheses can actually hold.

II.159.3 Additive robustness to corrupted domain points

Theorem II.159.5 (Robust shared interpolation). *Suppose $B^\circ \subseteq B \subseteq H \times J$ has weighted complement within B of normalized mass at most δ :*

$$\frac{1}{mW} \sum_h w_h |B_h \setminus B_h^\circ| \leq \delta.$$

Assume that f restricted to each B_h° can be partitioned into at most q polynomial classes of degree at most $t-1$. Then there is a shared r -subset R and a family of the forms (II.159.2) and (II.159.3), using the actual sections of f on B , which covers all but at most

$$\delta + \frac{tq(m-r)}{m(r+1)} \quad (\text{II.159.5})$$

of the ambient mass in B . The exact class-profile bound may replace the second term. No factor of r multiplies δ .

Proof. Apply Theorem II.153.2 to the good classes inside B° . The points in B° that fail certification have at most the stated sampling mass. Every certified nonsampled good point has t anchors in its own good class, where the actual values of f agree with the class polynomial. Consequently the candidate built from those actual sections is correct. A good sampled point is covered directly.

Candidates using a corrupted anchor can be ignored for this correctness argument: their presence only adds possible graphs. Discarding the corrupt points themselves costs at most δ , and discarding uncertified good points costs at most the sampling term. The losses add. $\square$

The theorem does not claim that the unknown good set or its class partition can be discovered efficiently from corrupted data. It is an existence and list-cover result. The conditional-expectation algorithm applies directly when the good class descriptions are supplied.

II.160 Labelled slice covers and candidate counts

Source 39, delivered Section 10.

The sampling theorem controls how many horizontal sections are needed. A second, separate issue is how the candidate lists for those sections are carried through a common refinement. Replacing each individual slice list by one large union list is safe but can be very wasteful.

Let the base H have a partition into cells U . On each cell let $\mathcal{V}_U \subseteq \mathbb{F}^U$ be a linear space. Suppose that for each sampled anchor $a \in R$ there is a list

$$\mu_{U,a,1}, \dots, \mu_{U,a,L_a} \in \mathcal{V}_U$$

covering the partial section ψ_a at every good base point of its domain. The list sizes may depend on U ; uniform upper bounds are used below to avoid clutter. The phrase “covering” means that at least one member equals $\psi_a(h)$ at the point in question. No globally consistent branch is assumed.

Theorem II.160.1 (Labelled interpolation of slice lists). *Let $H_0 \subseteq H$ be a common good base set for the slice lists above. On $U \times J$, the points certified by Theorem II.159.1 whose base lies in H_0 are covered by at most*

$$\sum_{a \in R} L_a + \sum_{\substack{S \subseteq R \\ |S|=t}} \prod_{a \in S} L_a \quad (\text{II.160.1})$$

candidates in

$$\mathcal{V}_U \otimes \text{Poly}_{<t}(J) := \left\{ \sum_{j=0}^{t-1} v_j(h) x^j : v_j \in \mathcal{V}_U \right\}.$$

For a common bound $L_a \leq L$, the count is

$$rL + \binom{r}{t} L^t. \quad (\text{II.160.2})$$

If the base exceptional set has weighted mass at most ϵW , the total ambient exceptional mass is at most the sampling loss plus ϵ .

Proof. For each anchor a and each list member, include the direct lift $\mu_{U,a,j}(h)$, constant in x . For each t -subset $S \subseteq R$ and each choice of a list index j_a for every $a \in S$, include

$$\nu_{U,S,(j_a)}(h, x) = \sum_{a \in S} L_{S,a}(x) \mu_{U,a,j_a}(h). \quad (\text{II.160.3})$$

The direct lifts belong to the stated space because constants have degree less than t . Expanding the coefficients $L_{S,a}(x)$ in powers of x shows that every function in (II.160.3) also belongs to that space, using linearity of $\mathcal{V}_U$.

If a certified point is directly sampled, select a correct member of that anchor’s list. Otherwise choose its t valid anchors from the same class. At its base point $h \in H_0$, each of those t anchor sections has at least one correct list member. Select these members individually. Formula (II.159.3) shows that the resulting candidate is correct at the target. A choice for other anchors is irrelevant.

The number of direct candidates is $\sum_a L_a$. For a fixed S , the number of possible assignments is $\prod_{a \in S} L_a$, giving (II.160.1). Removing a base set of weight ϵW deletes at most $m\epsilon W$ ambient mass, so the exceptional fractions add. $\square$

[write] 7 October 2026; sources 60–69. *The direct-anchor term is unnecessary* (source 67, Section II.338): the t -anchor interpolants also cover the sampled points, so $E_t(L_1, \dots, L_r)$ candidates suffice, exactly in this interface (its Theorems II.339.1 and II.339.6); the affine case is formalized since 081abc4d3 (Section16SampledOrAnchoredOn.compressed_multilinear_cover, Proofs16CompressedInterpolation.lean, line 56, at HEAD fb9602e55).

Corollary II.160.2 (Affine and multilinear specialization). *For affine fibres, the bound is*

$$rL + \binom{r}{2} L^2. \quad (\text{II.160.4})$$

If $\mathcal{V}_U$ consists of functions affine separately in every base coordinate, all the lifted candidates are affine separately in every base coordinate and in the final coordinate x .

Proof. Set $t = 2$ in Theorem II.160.1. Each coefficient in (II.159.4) is affine in x and independent of the base variables, while each list member is independent of x . The sum has degree at most one in each variable separately. $\square$

The mixed list count in (II.160.1) is the sum of the first and t th elementary symmetric polynomials in the list sizes. This formulation is useful when the horizontal sections have unequal complexity. The uniform bound (II.160.2) is a convenient corollary, not the sharpest piece of information furnished by the construction.

There is no claim that all these expressions are different. The count is an upper bound on a concrete generated list, not a matching lower bound on the number of distinct necessary graphs. Algebraic coincidences, zero coefficients, repeated lists, or extra relations can reduce the number.

II.161 Hereditary refinement without loss of slice labels

Source 39, delivered Section 11.

A common refinement is only useful if it preserves the required geometric scale. The next theorem states an explicit abstract hypothesis that is sufficient. It does not assume that arbitrary intersections of progression partitions retain large scale.

Definition II.161.1 (Hereditary single-slice refinement). Let a finite base set H carry an additive nonnegative measure μ with $\mu(H) > 0$. Suppose there is a class of admissible nonempty cells, a scale $\text{sc}(U) \geq 1$ for each such cell, and linear candidate spaces $\mathcal{V}_U \subseteq \mathbb{F}^U$. Assume restrictions of members of $\mathcal{V}_U$ to an admissible subcell $U' \subseteq U$ belong to $\mathcal{V}_{U'}$.

Let $\psi_1, \dots, \psi_r$ be partial functions on H . Assume that for each i , every admissible U , and each error parameter τ under consideration, there is a finite partition of U into admissible cells U' such that

$$\text{sc}(U') \geq \text{sc}(U)^{a(\tau)}, \quad 0 < a(\tau) \leq 1, \quad (\text{II.161.1})$$

and an exceptional set in U of measure at most $\tau\mu(U)$, outside which ψ_i on each U' is covered by a list of at most $L(\tau)$ members of $\mathcal{V}_{U'}$. The same functions a, L are valid for every i and every cell encountered in the refinement.

If a concrete refinement lemma has additional lower-size hypotheses, those hypotheses must hold at every stage where it is invoked. They are not removed by Definition II.161.1. In an application, admissibility can encode properness, coordinate restrictions, or inheritance from a specified parent progression.

Theorem II.161.2 (Label-preserving common refinement). *Suppose H itself is admissible, with scale M . Under Definition II.161.1, choose errors $\tau_1, \dots, \tau_r > 0$. There is a common admissible partition of H and a common good base set such that the total exceptional measure is at most $\mu(H) \sum_i \tau_i$, every final cell has scale at least*

$$M^{\prod_{i=1}^r a(\tau_i)}, \quad (\text{II.161.2})$$

and on each final cell, slice i retains its own list of at most $L(\tau_i)$ candidates. It is not necessary to replace that list size by $\prod_i L(\tau_i)$.

Proof. Refine the initial cell using the single-slice property for ψ_1 at error τ_1 . Then refine every current cell using the property for ψ_2 at error τ_2 , and continue. At each stage retain, by restriction, the lists already attached to all earlier slices. Restriction closure ensures that these remain valid candidates and does not increase their sizes.

At stage i , the new exceptional sets in the parent cells have total measure at most $\tau_i \sum_U \mu(U) = \tau_i \mu(H)$ because the parent cells partition H . The union of all stage exceptions therefore has measure at most $\mu(H) \sum_i \tau_i$.

For scales, induction gives after stage i the lower bound $M^{\prod_{j=1}^i a(\tau_j)}$: if the parent scale is at least the previous bound, then raising it to the positive power $a(\tau_i)$ preserves the inequality. Once a slice acquires its list at its own stage, later stages only restrict it. Thus its final list has at most $L(\tau_i)$ elements, as claimed. $\square$

Corollary II.161.3 (Equal error allocation). *For $\tau_i = \epsilon/r$, the common good set loses at most ϵ of the base mass, the scale bound is*

$$M^{a(\epsilon/r)^r}, \quad (\text{II.161.3})$$

and each slice list has size at most $L(\epsilon/r)$. Combined with affine shared-anchor selection, the graph count per base cell is

$$rL(\epsilon/r) + \binom{r}{2} L(\epsilon/r)^2. \quad (\text{II.161.4})$$

The total ambient exceptional fraction is at most

$$\frac{2q(m-r)}{m(r+1)} + \epsilon. \quad (\text{II.161.5})$$

Proof. Apply Theorem II.161.2, then Theorem II.160.1 and Theorem II.153.2 with $t = 2$. All interpolation anchors lie in the same row as their target, so restricting to a common good base point preserves all the required anchor lists. $\square$

II.161.1 Unequal slice costs and a parameter frontier

If slice i has its own functions $a_i(\tau_i)$ and $L_i(\tau_i)$, the same proof gives scale $M^{\prod_i a_i(\tau_i)}$ and a labelled candidate bound

$$\sum_i L_i(\tau_i) + \sum_{|S|=t} \prod_{i \in S} L_i(\tau_i), \quad \sum_i \tau_i \leq \epsilon. \quad (\text{II.161.6})$$

This gives an explicit optimization problem for error allocation; equal errors need not be optimal for heterogeneous slices. The theorem asserts the whole feasible frontier (II.161.6). It does not solve an unspecified optimization without information about the cost functions.

The asymmetric version can also use the exact profile loss (II.153.2) in place of the universal sampling term. Thus the method retains three separate forms of information: the row class-size profile, the number of anchors actually chosen, and the list cost of each horizontal section. Collapsing all of them to a single worst-case parameter is optional rather than built into the proof.

II.162 Source 34's Appendix A: the sharp nonuniform transport threshold [write]

Added in the seventh write, 6 October 2026.

The next section is source 34's Appendix A. *Logically it belongs in source 07's chapter on transport rigidity* (Section II.11), after its Corollary II.15.3; it is printed here because no number printed earlier may move. In source 07's finite-fibre directional model the sufficient threshold $1/(2[1 + (k-1)\kappa]) = 1/(2 + (2k-2)\kappa)$ is optimal for every Freiman order $k \geq 2$ and every fibre ratio $\kappa \geq 1$, already with zero boundary loss, so a boundary bound alone gives no further improvement (Theorem II.163.1). This answers source 07's Question 1 (item II.9); source 07 had sharpness at $\kappa = 1$ (Theorem II.16.1). The intake checked the sharpness family for four parameter triples, including the printed example on $\mathbb{Z}/8\mathbb{Z}$. The formal `lemma_10_9` (`Section10.lean`, line 335; companion `lemma_10_9_holds`, `Proofs10Shift.lean`, line 717, at HEAD 89c1cbc38) concerns the prime-order interface and is unaffected, as source 34 says. Source 34's introduction is Section F.39.

II.163 The sharp nonuniform directional rigidity threshold

Source 34, delivered Appendix A.

This appendix supplies an independent local result for the inverse part of the argument. Source 07 of the merged report proves a sufficient nonuniform directional threshold and asks whether it is sharp [155]. The following construction settles that question in the same finite-fiber model. The sharpness examples have no boundary loss.

II.163.1 The precise sampling model

Let G, H be finite abelian groups, let $S \subseteq G$ be nonempty, and let X_g be a finite nonempty fiber for $g \in S$. Write

$$R_g = |X_g|, \quad \Omega = \sum_{g \in S} R_g, \quad \mu(g) = R_g/\Omega, \quad \frac{\max_g R_g}{\min_g R_g} \leq \kappa.$$

Let $\phi : \bigsqcup_{g \in S} X_g \rightarrow H$, let $D \subseteq G$, and let $\psi : D \rightarrow H$. For a direction $d \in D$, draw $g \sim \mu$ and x uniformly from X_g . If $g + d \notin S$, declare failure. Otherwise draw y independently and uniformly from X_{g+d} and test

$$\phi(y) - \phi(x) = \psi(d).$$

Let δ_d be its failure probability. The map ψ is a Freiman homomorphism of order k if every relation $d_1 + \cdots + d_k = e_1 + \cdots + e_k$ with directions in D , repetitions allowed, implies equality of the corresponding image sums.

Theorem II.163.1 (The exact universal strict threshold). *For every integer $k \geq 2$ and every real $\kappa \geq 1$, the optimal universal strict threshold is*

$$\delta_{\text{crit}}(k, \kappa) = \frac{1}{2 + (2k - 2)\kappa}. \quad (\text{II.163.1})$$

If every $\delta_d < \delta_{\text{crit}}(k, \kappa)$, then ψ is a Freiman homomorphism of order k . No larger universal strict threshold is valid. For rational κ , a counterexample attains equality in (II.163.1). Sharpness holds with $S = G$, hence with zero exit probability in every direction.

Proof. We first recall the sufficient bound with its sampling weights. Fix a relation $\sum_i d_i = \sum_i e_i$. Choose one common initial fiber with index $g \sim \mu$. Follow the two paths of partial sums. Where their indices are in S , choose uniform fiber samples, using the same initial and terminal samples for both paths and independent internal samples. Samples at distinct path positions are drawn separately even if their fiber indices coincide. Leaving S is a failure.

Each first edge has its prescribed defect at most $\delta := \max_{d \in D} \delta_d$. For a later edge, its starting index is $g + s$ for a fixed partial sum s . On indices for which both g and $g + s$ are in S ,

$$\mu(g) \leq \kappa \mu(g + s).$$

Therefore the probability of failure at that edge, restricted to paths that have not already left S , is at most $\kappa\delta$. This includes the first exit at that edge. A union bound over the two paths gives total failure probability at most

$$[2 + (2k - 2)\kappa]\delta < 1.$$

Some realization has every edge correct. Telescoping along its two paths, with their common endpoints, proves the required image relation. For an empty D the assertion is immediate; otherwise the maximum is attained because G is finite.

For sharpness, fix an integer $m \geq 2$ and positive integers $a \leq b$. Put $N = 2km$ and $G = H = \mathbb{Z}/N\mathbb{Z}$, with $S = G$. Partition the representatives $0 \leq r < N$ into $2k$ consecutive blocks of length m . Give every point-index in blocks 0 and $2k - 1$ a fiber of size a , and every other point-index a fiber of size b . On the fiber indexed by r , let ϕ be constant with value $\lfloor r/m \rfloor \in \mathbb{Z}/N\mathbb{Z}$. Take

$$D = \{m, -m\}, \quad \psi(m) = 1, \quad \psi(-m) = -1.$$

Only the final block fails in direction m , and only the first block fails in direction $-m$. Consequently

$$\delta_m = \delta_{-m} = \frac{a}{2a + (2k - 2)b}. \quad (\text{II.163.2})$$

There is no boundary failure. Yet the domain relation consisting of k copies of m and k copies of $-m$ is valid modulo $N = 2km$, whereas its image relation $k = -k$ fails because $N > 2k$. The map is not Freiman of order k .

For rational κ , take $b/a = \kappa$ to attain the stated threshold. For arbitrary real κ , choose rational ratios $b/a \leq \kappa$ tending to κ . The defects in (II.163.2) tend to (II.163.1) from above, so every larger proposed universal strict threshold admits a counterexample. $\square$

Example II.163.2. At $k = 2$, $\kappa = 3$, take $m = 2$, $a = 1$, and $b = 3$. The group is $\mathbb{Z}/8\mathbb{Z}$, the fiber sizes by block are $(1, 3, 3, 1)$, and each block has two indices. Both directional defects are $1/8$. The relation $2 + 2 = (-2) + (-2)$ holds modulo eight, but $1 + 1 \neq (-1) + (-1)$. This realizes equality at the threshold $1/(2 + 2 \cdot 3)$.

Corollary II.163.3 (A boundary bound alone gives no further improvement). *Within this model, adding an arbitrary upper bound on the exit probabilities, including the bound zero, does not increase the universal threshold in Theorem II.163.1.*

Proof. Every sharpness example in the theorem has $S = G$ and exit probability zero. $\square$

The sufficient threshold is credited to source 07; the matching nonuniform example resolves its sharpness question. The examples use composite cyclic groups. They do not determine the maximal Freiman order available under Gowers’s original prime-order and nested-set hypotheses at Lemma 10.9. Those additional hypotheses can carry information absent from this general model.

II.164 Source 40’s chapter: one-stage Fejér selection at Lemma 15.5 [write]

Added in the eighth write, 7 October 2026.

The next eight sections are source 40’s Sections 2–9, on the arrangement-selection step of Gowers’s Lemma 15.5. *Logically the chapter comes before source 35’s degree-weighted selection* (Section II.144), which builds on it and cites it ahead of its placement, and after source 17’s Fejér restriction (its Section 9) and the Section 15 material of sources 25 and 27; it is printed here because no number printed earlier may move. Source 40’s introduction, with its explicit theorem, is Section F.45; its checks and questions are in Part V.

What is proved. For ordered d -arrangements of k -cubes in a density- β set $B \subseteq \mathbb{F}_q^{k+1}$, a single normalized Fejér kernel of bandwidth $L = \lceil 2(s+1)/(\alpha\eta) \rceil$, $s = 2d2^k$, applied to the single feature $F(y, r) = r \prod_i y_i$, retains at least $\alpha^s \eta^{s-1} \beta^{2d-1} q^M / (3^{s-1} (s+1)^s)$ arrangements, a proportion $1/(1+\eta)$ of them respected, as soon as $p > 2(L-1)$ and $q \geq (k+1)\beta^{-(2d-1)}(18(s+1)^2/(\alpha\eta)^2)^s$ (Theorem F.45.1; the exact form is Theorem II.169.2). Gowers’s repeated degree-one Riesz products lose the exponent $s2^{s-1}$; here the exponents are s and $s-1$, and the sufficient field size is also smaller (Corollary II.170.1 at $d=8$). The feature has exactly the intended universal relation space (Theorem II.166.2), and all other bounded-coefficient relations are counted by a zero-count with repeated vertices treated explicitly (Theorem II.167.3). Vector-valued targets cost nothing, target torsion is computed exactly (Corollary II.171.1, Proposition II.171.2), and within homogeneous positive-Fourier kernels a gain Q forces good survival at most $Q^{-(s-1)}$, attained by Fejér kernels (Theorem II.172.2, Proposition II.172.3). Source 35 improves the field-size exponent from $2s$ to $E_{k,m} \sim 2s/k$ and the retention constant to $\alpha^s \eta^{s-1} / (2(s+2)^s)$ (note after Theorem II.169.2).

Relation to the formal statements (checked at HEAD 6d4942222). `lemma_15_5` (Sections14_15.lean, line 414) has the exact companion `lemma_15_5_holds` (Proofs15Restriction.lean, line 3116), with an existential threshold. Whether Corollary II.170.1 gives it with an explicit threshold depends on matching the formal count conventions, as for source 35 (marked question, item II.50). Nothing in this chapter is formalized.

II.165 Arrangements, respect, and the counting normalization

Source 40, delivered Section 2.

Throughout, $\mathbb{F}_q$ denotes a finite field. Write $E_k = \{0, 1\}^k$, $\pi(\varepsilon) = (-1)^{|\varepsilon|}$, and

$$\sigma_j = \begin{cases} 1 & 1 \leq j \leq d, \\ -1 & d < j \leq 2d. \end{cases}$$

Products and sums occurring in coordinates are taken in $\mathbb{F}_q$. The symbols π and σ also denote their images in that field.

Definition II.165.1 (Ordered d -arrangement). An ordered d -arrangement is a choice of

$$h \in \mathbb{F}_q^k, \quad x^{(1)}, \dots, x^{(2d)} \in \mathbb{F}_q^k, \quad r_1, \dots, r_{2d} \in \mathbb{F}_q, \quad \sum_{j=1}^{2d} \sigma_j r_j = 0.$$

Its labeled vertices are

$$v_{j,\varepsilon} = (x^{(j)} + \varepsilon \odot h, r_j), \quad 1 \leq j \leq 2d, \quad \varepsilon \in E_k, \quad (\text{II.165.1})$$

where $\odot$ denotes coordinatewise multiplication. It is an arrangement in B when every vertex belongs to B . Set

$$s = 2d2^k, \quad M = k + 2dk + (2d - 1) = (2d + 1)k + 2d - 1, \quad (\eta_0)_{j,\varepsilon} = \sigma_j \pi(\varepsilon).$$

For $\phi : B \rightarrow \mathbb{F}_q$, its defect on the arrangement is

$$\Delta_\phi(\boldsymbol{\rho}) = \sum_{j,\varepsilon} \sigma_j \pi(\varepsilon) \phi(v_{j,\varepsilon}). \quad (\text{II.165.2})$$

The arrangement is *respected* when this defect is zero. Let $T(B)$ be the number of all arrangements in B and $G_\phi(B)$ the number of respected ones.

There are exactly q^M arrangements in the ambient space: the last r coordinate is uniquely determined by the first $2d - 1$. Labeling is important. We never divide by a cube symmetry group, and we do not remove coincident configurations before forming $T(B)$.

Lemma II.165.2 (Density-sensitive total count). For $B \subseteq \mathbb{F}_q^{k+1}$ of density β ,

$$T(B) \leq T_0 := \beta^{2d-1} q^M. \quad (\text{II.165.3})$$

Proof. For fixed h and r , let $C_h(r)$ be the number of anchors $x \in \mathbb{F}_q^k$ for which the whole cube $(x + \varepsilon \odot h, r)$ lies in B . Necessarily

$$0 \leq C_h(r) \leq q^k, \quad \sum_{r \in \mathbb{F}_q} C_h(r) \leq |B|,$$

since the anchor itself is a vertex. Once $r_1, \dots, r_{2d-1}$ are fixed, r_{2d} is determined. Hence the arrangement count at this h is at most

$$q^k \sum_{r_1, \dots, r_{2d-1}} \prod_{j < 2d} C_h(r_j) = q^k \left(\sum_r C_h(r) \right)^{2d-1} \leq q^k |B|^{2d-1}.$$

Summing over the q^k possible h gives $q^{2k} |B|^{2d-1} = \beta^{2d-1} q^M$, as claimed. $\square$

Remark II.165.3 (Why this normalization is natural). For $B = \mathbb{F}_q^{k+1}$ there is equality. More generally, if $B = \mathbb{F}_q^k \times R$ for an additive subgroup $R \leq (\mathbb{F}_q, +)$, then all the final r coordinates determined by the constraint remain in R , and there is again equality in (II.165.3). Thus the factor β^{2d-1} cannot be universally reduced for finite fields. This observation concerns the total-count normalization, not optimality of the later selection theorem.

II.166 A single monomial isolates the universal relation

Source 40, delivered Section 3.

Define

$$P(y) = \prod_{i=1}^k y_i, \quad F(y, r) = rP(y). \quad (\text{II.166.1})$$

The next lemma strengthens the hypothesis in the cube-moment step underlying Gowers' Corollary 15.2: constancy of the top monomial alone is enough. No separate hypotheses for every squarefree submonomial are needed.

Lemma II.166.1 (Top-monomial cube rigidity). *Let K be any field, let each $h_i \in K$ be nonzero, and let $\eta : E_k \rightarrow K$. The formal polynomial*

$$U_\eta(x) = \sum_{\varepsilon \in E_k} \eta_\varepsilon \prod_{i=1}^k (x_i + \varepsilon_i h_i)$$

is constant in $x_1, \dots, x_k$ if and only if $\eta_\varepsilon = c\pi(\varepsilon)$ for some $c \in K$. In that case

$$U_\eta(x) = c(-1)^k \prod_i h_i. \quad (\text{II.166.2})$$

The assertion is valid in every characteristic.

Proof. For $A \subseteq [k]$, set $M_A = \sum_\varepsilon \eta_\varepsilon \prod_{i \in A} \varepsilon_i$. Expansion gives

$$U_\eta(x) = \sum_{A \subseteq [k]} M_A \prod_{i \in A} h_i \prod_{i \notin A} x_i.$$

The distinct x monomials are linearly independent. Since the h_i are nonzero, constancy is equivalent to $M_A = 0$ for every proper A . Writing $S = \text{supp}(\varepsilon)$, the moments are upper subset sums $M_A = \sum_{S \supseteq A} \eta_S$. Finite inclusion–exclusion gives

$$\eta_S = \sum_{A \supseteq S} (-1)^{|A|-|S|} M_A.$$

Indeed, after substituting for M_A , the coefficient of η_T is $\sum_{S \subseteq A \subseteq T} (-1)^{|A|-|S|}$, which is one for $T = S$ and zero otherwise. Thus $\eta_S = (-1)^{k-|S|} M_{[k]} = c(-1)^{|S|}$ with $c = (-1)^k M_{[k]}$. The constant term is $M_{[k]} \prod_i h_i$, giving (II.166.2). The reverse implication follows from the same expansion, or from $\prod_i (x_i - (x_i + h_i)) = (-1)^k \prod_i h_i$. $\square$

Eliminate r_{2d} using the arrangement equation:

$$r_{2d} = \sum_{j < 2d} \sigma_j r_j. \quad (\text{II.166.3})$$

The free parameters then form K^M .

Theorem II.166.2 (Universal relation space for one feature). *For coefficients $\eta = (\eta_{j,\varepsilon}) \in K^s$, the formal polynomial*

$$\mathcal{P}_\eta(h, x, r) = \sum_{j,\varepsilon} \eta_{j,\varepsilon} F(v_{j,\varepsilon}), \quad (\text{II.166.4})$$

with (II.166.3) substituted, is identically zero if and only if $\eta \in K\eta_0$. If it is not zero, its total degree is at most $k + 1$.

Proof. Put

$$U_j(x^{(j)}, h) = \sum_{\varepsilon} \eta_{j,\varepsilon} P(x^{(j)} + \varepsilon \odot h).$$

After substitution, the polynomial in (II.166.4) is

$$\sum_{j < 2d} r_j (U_j + \sigma_j U_{2d}). \quad (\text{II.166.5})$$

Suppose it is identically zero. Specialize all h_i to one. Each coefficient $U_j + \sigma_j U_{2d}$ of the independent variable r_j is zero. The two summands depend on disjoint anchor variables. Consequently both are constant in their own anchor variables: this follows, for example, by fixing the other anchor at zero and comparing polynomial coefficients. As j ranges over $1, \dots, 2d - 1$, every block is covered. Lemma II.166.1 now gives $\eta_{j,\varepsilon} = c_j \pi(\varepsilon)$ for each j .

For general h the coefficient of r_j is therefore $(-1)^k (\prod_i h_i) (c_j + \sigma_j c_{2d})$. Its vanishing gives $c_j = -\sigma_j c_{2d}$, so with $c = -c_{2d}$ we have $c_j = c \sigma_j$ for all j , including $j = 2d$. Thus $\eta = c \eta_0$.

Conversely, if $\eta = c \eta_0$, the parity identity gives

$$\sum_{j,\varepsilon} \eta_{j,\varepsilon} F(v_{j,\varepsilon}) = c (-1)^k \left(\prod_i h_i \right) \sum_j \sigma_j r_j = 0.$$

Finally, each term is a product of $k + 1$ linear coordinate expressions in the free parameters, so its total degree is at most $k + 1$. $\square$

Remark II.166.3 (Formal polynomial versus finite-field function). The theorem is a polynomial identity statement over any field. It does not identify every pointwise relation in a small finite field with a formal identity. The passage to a large finite field is made through a zero-count in the next section. This distinction avoids a hidden small-characteristic assumption in the algebraic argument.

II.167 Bounded-coefficient exceptional arrangements

Source 40, delivered Section 4.

Let $H \geq 1$ and suppose $p > 2H$. Thus distinct integers in $[-H, H]$ have distinct images in $\mathbb{F}_q$.

Definition II.167.1 (H -regularity). An arrangement is H -regular if every vector $\eta \in \{-H, \dots, H\}^s$ satisfying

$$\sum_{j,\varepsilon} \eta_{j,\varepsilon} F(v_{j,\varepsilon}) = 0$$

is $a \eta_0$ for some integer $a \in [-H, H]$. All other arrangements are H -exceptional. Write D_H for the number of exceptional arrangements in the full ambient space.

This definition uses one feature, not the larger feature family in the source paper. Therefore it is a stronger regularity requirement. All signed multiples $a \eta_0$ really are relations by Theorem II.166.2.

Lemma II.167.2 (Polynomial zero-count). *A nonzero polynomial of total degree at most D in $m \geq 1$ variables over $\mathbb{F}_q$ has at most $D q^{m-1}$ zeros in $\mathbb{F}_q^m$.*

Proof. The one-variable case is the usual root bound. Induct on m . Write the polynomial as a polynomial of degree t in the last variable, with nonzero leading coefficient A of degree at most

$D - t$ in the other variables. There are at most $(D - t)q^{m-2}$ choices where A vanishes, and at most t last-coordinate roots for every other choice. Bounding the first group by q roots each and the second by tq^{m-1} in total yields at most Dq^{m-1} roots. The case of a nonzero constant leading coefficient is included, and the bound is harmless when it exceeds q^m . $\square$

Theorem II.167.3 (Single-feature exceptional count). *Set*

$$C_{k,s,H} = \frac{k+1}{2}((2H+1)^s - (2H+1)). \quad (\text{II.167.1})$$

Then

$$D_H \leq C_{k,s,H} q^{M-1} \leq (k+1)(2H+1)^s q^{M-1}. \quad (\text{II.167.2})$$

Every H -regular arrangement has s distinct physical vertices and, in particular, has all sidelengths nonzero.

Proof. The bounded coefficient box contains $(2H+1)^s$ integer vectors. Its intersection with the field line $\mathbb{F}_q \eta_0$ consists of exactly $2H+1$ vectors: one coordinate with sign $+1$ determines the scalar, and injectivity of reduction modulo p on $[-H, H]$ determines every other coordinate as the corresponding signed integer. These are just the intended multiples.

For each remaining coefficient vector, Theorem II.166.2 gives a nonzero polynomial of degree at most $k+1$ in the M free parameters. Its zero set has size at most $(k+1)q^{M-1}$ by Lemma II.167.2. The vectors η and $-\eta$ have the same zero set. None of the remaining vectors is zero, and their sign pairs have size two, giving exactly $((2H+1)^s - (2H+1))/2$ pairs. A union bound proves the first inequality in (II.167.2); the second only simplifies it.

If two labeled vertices v_a, v_b coincide, the vector with entries 1 at a , -1 at b , and zero elsewhere is a bounded relation. It is not a multiple of η_0 : it is nonzero, while $s \geq 4$ and a nonzero multiple of η_0 has no zero coordinate. Such an arrangement is exceptional. A zero sidelength forces precisely such a repeated vertex, proving the final assertion. $\square$

Remark II.167.4 (What the exceptional bound pays for). The bound includes zero sidelengths, repeated vertices, and all other bounded relations. No additional diagonal error is left to be inserted later. It is deliberately a union bound, not a claim that the corresponding algebraic hypersurfaces are disjoint or independent. The price of increasing H is explicit: roughly $(2H+1)^s/q$ of the ambient parameter space may have to be discarded from the exact survival calculation.

II.168 Fejér kernels and exact survival probabilities

Source 40, delivered Section 5.

II.168.1 The normalized kernel

For $L \geq 2$ define a function on the unit circle by

$$K_L(z) = \left| \frac{1}{L} \sum_{n=0}^{L-1} z^n \right|^2 = \sum_{a=-(L-1)}^{L-1} c_a z^a, \quad c_a = \frac{L - |a|}{L^2}. \quad (\text{II.168.1})$$

The coefficient formula follows by counting pairs (n, m) with $n - m = a$. All coefficients are nonnegative, $c_{-a} = c_a$, and $0 \leq K_L(z) \leq 1$ by the triangle inequality. Also $c_0 = 1/L$ and $\sum_a c_a = K_L(1) = 1$.

For an arrangement with s vertices set

$$b_{s,L} = c_0^s = L^{-s}, \quad (\text{II.168.2})$$

$$a_{s,L} = \sum_{a=-(L-1)}^{L-1} c_a^s = L^{-s} R_s(L), \quad (\text{II.168.3})$$

$$R_s(L) = 1 + 2 \sum_{a=1}^{L-1} \left(1 - \frac{a}{L}\right)^s = 1 + 2L^{-s} \sum_{a=1}^{L-1} a^s. \quad (\text{II.168.4})$$

Lemma II.168.1 (Finite gain bounds). *For integers $s \geq 1$ and $L \geq 2$,*

$$\frac{2L}{s+1} \leq R_s(L) \leq L. \quad (\text{II.168.5})$$

For fixed s , as $L \rightarrow \infty$,

$$\frac{R_s(L)}{L} \rightarrow \frac{2}{s+1}, \quad a_{s,L} \sim \frac{2}{s+1} L^{1-s}. \quad (\text{II.168.6})$$

Proof. The trapezoidal sum for the convex function $f(x) = (1-x)^s$ on $[0, 1]$ lies above its integral. That sum is $R_s(L)/(2L)$, whereas the integral is $1/(s+1)$. For the upper bound, use $(1-a/L)^s \leq 1-a/L$ in (II.168.4), giving $R_s(L) \leq 1 + (L-1) = L$. Finally, the trapezoidal sums converge to the integral by continuity; multiplying by the powers of L in (II.168.3) gives (II.168.6). $\square$

II.168.2 The correlated random subset

Choose a nontrivial additive character $\chi : \mathbb{F}_q \rightarrow \mathbb{C}$ with $|\chi(x)| = 1$. For example, if $q = p^m$, take $\chi(x) = \exp(2\pi i \operatorname{Tr}_{\mathbb{F}_q/\mathbb{F}_p}(x)/p)$, interpreting the trace as an integer modulo p . The elementary orthogonality relation is

$$\mathbb{E}_{u \in \mathbb{F}_q} \chi(uz) = \mathbf{1}_{z=0}. \quad (\text{II.168.7})$$

For $z \neq 0$, multiplication by z permutes the field and the sum of a nontrivial character is zero: translating that sum by an element where the character is not one multiplies it by a scalar factor different from one while leaving it unchanged.

Choose $t, u \in \mathbb{F}_q$ uniformly and independently. Conditional on these two seeds, select each physical point $v \in B$ independently with probability

$$w_{t,u}(v) = K_L(\chi(t\phi(v) + uF(v))). \quad (\text{II.168.8})$$

Let B' be the resulting random subset. Only two field-valued seeds are used for the correlated part of the selection; the pointwise Bernoulli choices are additional randomness.

Lemma II.168.2 (Exact labeled moment). *For any labeled arrangement in B ,*

$$\begin{aligned} & \mathbb{E}_{t,u} \prod_{j,\varepsilon} w_{t,u}(v_{j,\varepsilon}) \\ &= \sum_{\eta \in [-(L-1), L-1]^s \cap \mathbb{Z}^s} \left(\prod_{j,\varepsilon} c_{\eta_{j,\varepsilon}} \right) \mathbf{1}_{\sum \eta_{j,\varepsilon} F(v_{j,\varepsilon}) = 0} \mathbf{1}_{\sum \eta_{j,\varepsilon} \phi(v_{j,\varepsilon}) = 0}. \end{aligned} \quad (\text{II.168.9})$$

In particular, every summand is nonnegative.

Proof. Expand each factor using (II.168.1). Character multiplicativity turns the product into $\chi(u \sum \eta F + t \sum \eta \phi)$ times its coefficient. Averaging independently over u and t applies (II.168.7) twice. All sums are finite, so their interchange needs no limiting argument. $\square$

Proposition II.168.3 (Good and bad survival). *Assume $p > 2(L - 1)$ and write $H = L - 1$. Every respected arrangement, whether regular or exceptional, survives in B' with probability at least $a_{s,L}$. Every H -regular nonrespected arrangement survives with probability exactly $b_{s,L}$. Every H -regular respected arrangement survives with probability exactly $a_{s,L}$.*

Proof. First consider the labeled moment in (II.168.9). For a respected arrangement, all intended coefficient vectors $\eta = a\eta_0$ contribute. They satisfy the feature equation by Theorem II.166.2, and their ϕ sum is $a\Delta_\phi = 0$. Since $c_{-a} = c_a$, their total contribution is $\sum_a c_a^s = a_{s,L}$. All other contributions are nonnegative.

There is a distinction between the labeled moment and actual subset survival when vertices repeat. Conditional on the seeds, actual survival is the product of $w_{t,u}(v)$ over *distinct* physical vertices. As each weight lies in $[0, 1]$, this product is at least the product over all labeled occurrences. Averaging therefore preserves the lower bound $a_{s,L}$, even for exceptional respected arrangements.

For an H -regular arrangement the vertices are distinct by Theorem II.167.3, so the labeled moment is exactly its survival probability. Regularity allows only $\eta = a\eta_0$ in that moment. If $\Delta_\phi \neq 0$, the remaining equation is $a\Delta_\phi = 0$ in $\mathbb{F}_q$. The only $a \in [-H, H]$ satisfying it is $a = 0$, since $p > 2H$. The sole contribution is then $c_0^s = b_{s,L}$. If the defect is zero, all intended multiples contribute and the sum is exactly $a_{s,L}$. $\square$

Remark II.168.4 (Two easily missed issues). Nonnegativity of the Fourier coefficients is what permits us to retain only the intended terms for a lower bound; it is not merely a convenient normalization. Separately, conditional independence applies to physical points, not to vertex labels. The repeated-vertex inequality above is necessary before any first-moment estimate for the good arrangements is justified.

II.169 The one-stage selection theorem

Source 40, delivered Section 6.

II.169.1 An abstract first-moment step

The following elementary observation separates the probabilistic selection argument from the specific choice of kernel.

Lemma II.169.1 (Score selection). *Suppose a family of configurations has at most T_0 members, at least αT_0 of them good. A random subset retains each good configuration with probability at least $a > 0$. Each nonexceptional bad configuration is retained with probability at most b , and there are at most D exceptional bad configurations. If*

$$\alpha a - \frac{b + D/T_0}{\eta} > 0, \quad (\text{II.169.1})$$

there is an outcome with at least

$$\left(\alpha a - \frac{b + D/T_0}{\eta} \right) T_0 \quad (\text{II.169.2})$$

good configurations, in which the proportion of good configurations is at least $1/(1+\eta)$. In particular, if $D \leq bT_0$ and $a/b \geq 4/(\alpha\eta)$, this lower bound is at least $\alpha aT_0/2$.

Proof. Let X and Y count the retained good and bad configurations. By linearity of expectation, without any independence assumption between configurations,

$$\mathbb{E}X \geq \alpha aT_0, \quad \mathbb{E}Y \leq bT_0 + D.$$

Thus the real score $Z = X - Y/\eta$ has expectation at least (II.169.2). There is an outcome with $Z \geq \mathbb{E}Z$. For it, $X \geq Z$ and $Y \leq \eta X$, so $X/(X+Y) \geq 1/(1+\eta)$. The denominator is positive by (II.169.1). Under the final hypotheses, $(b + D/T_0)/\eta \leq 2b/\eta \leq \alpha a/2$. $\square$

II.169.2 The exact finite-bandwidth form

Theorem II.169.2 (Exact one-stage Fejér selection). *Let $k, d \geq 1$, $s = 2d^{2^k}$, and $M = (2d + 1)k + 2d - 1$. Let $B \subseteq \mathbb{F}_q^{k+1}$ have density $\beta > 0$, and suppose $\phi : B \rightarrow \mathbb{F}_q$ respects at least αT_0 arrangements, where $T_0 = \beta^{2d-1}q^M$ and $0 < \alpha, \eta \leq 1$. Choose an integer $L \geq 2$ for which*

$$p > 2(L-1), \quad R_s(L) \geq \frac{4}{\alpha\eta}, \quad \beta^{2d-1}q \geq C_{k,s,L-1}L^s. \quad (\text{II.169.3})$$

There is $B' \subseteq B$ for which

$$G_\phi(B') \geq \frac{\alpha}{2}L^{-s}R_s(L)T_0, \quad \frac{G_\phi(B')}{T(B')} \geq \frac{1}{1+\eta}. \quad (\text{II.169.4})$$

More generally, for any L satisfying only the characteristic condition, the same construction gives a positive lower bound

$$G_\phi(B') \geq \left(\alpha L^{-s}R_s(L) - \frac{L^{-s}}{\eta} - \frac{C_{k,s,L-1}}{\eta\beta^{2d-1}q} \right) T_0 \quad (\text{II.169.5})$$

with the same good proportion whenever the displayed coefficient is strictly positive.

Proof. Use the random subset in (II.168.8). Lemma II.165.2 bounds the total family by T_0 ; Proposition II.168.3 supplies the good probability $a = L^{-s}R_s(L)$ and the nonexceptional bad probability $b = L^{-s}$. The exceptional count is at most $C_{k,s,L-1}q^{M-1}$ by Theorem II.167.3. The last inequality in (II.169.3) makes this at most bT_0 , and the middle inequality gives the gain needed in Lemma II.169.1. This proves (II.169.4). Using the full score lower bound instead gives (II.169.5). $\square$

Proof of Theorem F.45.1. Take L from (F.45.1). Lemma II.168.1 gives $R_s(L) \geq 2L/(s+1) \geq 4/(\alpha\eta)$. As $\alpha\eta \leq 1$,

$$L \leq \frac{2(s+1)}{\alpha\eta} + 1 \leq \frac{3(s+1)}{\alpha\eta}. \quad (\text{II.169.6})$$

Also

$$\begin{aligned} C_{k,s,L-1}L^s &\leq (k+1)[L(2L-1)]^s \\ &\leq (k+1)(2L^2)^s \\ &\leq (k+1) \left(\frac{18(s+1)^2}{(\alpha\eta)^2} \right)^s. \end{aligned}$$

Thus the field-size hypothesis implies the last condition in (II.169.3). Theorem II.169.2 yields a subset with

$$\begin{aligned} G_\phi(B') &\geq \frac{\alpha}{2} L^{-s} R_s(L) T_0 \\ &\geq \frac{\alpha}{s+1} L^{1-s} T_0 \\ &\geq \frac{\alpha}{s+1} \left(\frac{\alpha\eta}{3(s+1)} \right)^{s-1} T_0 = \Gamma_s(\alpha, \eta) T_0. \end{aligned}$$

Since $T(B') \geq G_\phi(B')$, this proves the arrangement lower bound. The good proportion is at least $1/(1+\eta)$, which is at least $1-\eta$ because their difference is $\eta^2/(1+\eta)$.

Write $\beta' = |B'|/q^{k+1}$. Applying Lemma II.165.2 to B' gives $\Gamma_s \beta'^{2d-1} q^M \leq T(B') \leq (\beta')^{2d-1} q^M$. Taking positive $(2d-1)$ st roots proves (F.45.4).

For prime $q = p$, put $X = (s+1)/(\alpha\eta) \geq 5$. The size bound implies $q \geq (18X^2)^s \geq 6X \geq 2L$ by (II.169.6), and hence $p > 2(L-1)$. $\square$

[write] 7 October 2026; sources 40–46. *Improved by source 35* (written the same day, placed earlier, printed in Section II.144). Source 35's Theorem II.145.2 keeps the one-stage selector and the powers $\alpha^s \eta^{s-1}$ of this theorem, stratifies the coefficient arrays by degree, and lowers the field-size exponent from $2s$ to $E_{k,m} \sim 2s/k$ (23 against 128 at $d = 8, k = 2$), with retention $\alpha^s \eta^{s-1}/(2(s+2)^s)$, larger than $\alpha^s \eta^{s-1}/(3^{s-1}(s+1)^s)$ here for every $s \geq 2$. The intake's checks of this section (nullity one of the evaluation matrix for five (k, d) , the inequality $2L/(s+1) \leq R_s(L) \leq L$ for $s \leq 128, L \leq 300$, the constant chain in 36 exact cases) are recorded in the dossier.

[write] 7 October 2026; sources 54–59. *Used for Theorem 13.12.* Source 57 reproves this selection for Gowers's eight-arrangements with the single feature xy (its Theorem II.289.1, Section II.287) and propagates it to a written proof of Theorem 13.12 with density $(\alpha/2)^{2^{41}}$; the repository's formal proof of `theorem_13_12` uses its own uniform Fejér purification along that route (`section13_fejer_purification_uniform`, `Proofs13FejerUniformPurification.lean`, line 10, at HEAD `fe7165a3c`).

Remark II.169.3 (Parameters at the endpoints). When $\alpha = 1$, every possible arrangement in B is already respected by Lemma II.165.2, so one may simply take $B' = B$. When $\eta = 1$, the requested lower proportion $1 - \eta$ is vacuous; our theorem still gives the stronger proportion at least one half. The displayed constants are uniform and are not optimized at either endpoint. If $G_\phi(B) = 0$, no positive α hypothesis holds.

II.170 Comparison with Gowers' Lemma 15.5

Source 40, delivered Section 7.

The source result concerns $d = 8$, so

$$s = 2^{k+4}, \quad M = 17k + 15, \quad E_k = 2^{2^{k+4}+k+3} = s^{2^{s-1}}. \quad (\text{II.170.1})$$

In the repository's edited statement [192, Section 15], the retained arrangement lower bound is

$$\left(\frac{\alpha\eta}{4} \right)^{E_k} \beta^{15} N^{17k+15}, \quad (\text{II.170.2})$$

with sufficient largeness condition

$$N \geq k 3^s \beta^{-15} \left(\frac{4}{\alpha\eta} \right)^{E_k}. \quad (\text{II.170.3})$$

The transcription explicitly notes that the printed statement did not quantify α and omitted a largeness condition used by its proof. We compare against the corrected, meaningful form, not the unqualified statement.

Corollary II.170.1 (Replacement for the eight-arrangement selection). *Under the setup of Gowers’ Lemma 15.5, with N prime and $0 < \alpha, \beta, \eta \leq 1$, it is enough to assume*

$$N \geq (k+1)\beta^{-15} \left(\frac{18(s+1)^2}{(\alpha\eta)^2} \right)^s, \quad s = 2^{k+4}. \quad (\text{II.170.4})$$

There is $B' \subseteq B$ containing at least

$$\frac{\alpha^s \eta^{s-1}}{3^{s-1}(s+1)^s} \beta^{15} N^{17k+15} \quad (\text{II.170.5})$$

arrangements, with at least a proportion $1 - \eta$ respected.

Proof. Apply Theorem F.45.1 with $d = 8$ and use the final prime-field observation in its proof. $\square$

[write] **7 October 2026; sources 40–46.** *Against the formal statement* (checked at HEAD 6d4942222). `lemma_15_5` (Sections14_15.lean, line 414) has the exact companion `lemma_15_5_holds` (Proofs15Restriction.lean, line 3116), with an existential threshold and retention $(\alpha\eta/4)^{2^{2^{k+4}+k+3}} \beta^{15} N^{17k+15}$, the displayed (II.170.2). This corollary would give the formal conclusion with the explicit threshold (II.170.4), provided that source 40’s count of labelled arrangements, which includes repeated vertices and zero sidelengths, matches the conventions of `generalArrangementCount` and `respectedGeneralArrangementCount` (lines 229 and 234), and that the purity $1/(1+\eta)$ gives the formal form; that matching was not checked, so the observation joins source 35’s marked question (item II.50). The formal statement is already exact; nothing formal depends on it.

II.170.1 Both the loss and the sufficient field size improve

For $s \geq 32$ and $0 < \alpha, \eta \leq 1$, the lower bound (II.170.5) is larger than (II.170.2). To see this, $E_k \geq s$ and $E_k \geq s - 1$, so the ratio is bounded below by

$$\frac{4^{E_k}}{3^{s-1}(s+1)^s} > 1.$$

The last inequality follows, after taking s th roots, from $4^{2^{s-1}} > 3(s+1)$.

The sufficient threshold (II.170.4) is also smaller than (II.170.3). With $\delta = \alpha\eta \leq 1$, the ratio of the old threshold to the new one is

$$\frac{k}{k+1} \frac{4^{E_k}}{[6(s+1)^2]^s} \delta^{-(E_k-2s)}.$$

Here $E_k \geq 2s$, $k/(k+1) \geq 1/2$, and $4^{2^{s-1}} > 12(s+1)^2$ for $s \geq 32$. These facts make the ratio larger than one. Thus the result is not merely a better retained-mass bound purchased by a worse ambient-size assumption.

II.170.2 Why increasing bandwidth helps

At $L = 2$ the normalized Fejér kernel is the original probability kernel $(1 + \cos \theta)/2$, and

$$b_{s,2} = 2^{-s}, \quad R_s(2) = 1 + 2^{1-s}.$$

With r independent stages, regular good and bad survival probabilities are $a_{s,2}^r$ and $b_{s,2}^r$; the ratio is $R_s(2)^r$. Making it at least a prescribed $Q > 1$ takes

$$r \geq \frac{\log Q}{\log(1 + 2^{1-s})} \sim 2^{s-1} \log Q \quad \text{as } s \rightarrow \infty.$$

This explains the large exponent in the original route. In contrast, $R_s(L) \geq 2L/(s+1)$: a single stage with L proportional to $(s+1)Q$ produces the required ratio and has $a_{s,L}$ of order L^{1-s} for fixed s .

The trade-off has not disappeared. The higher kernel degree requires excluding a larger bounded coefficient box. The exceptional count pays for that box, and the explicit field-size hypothesis makes its contribution no larger than the ordinary bad survival mass.

II.170.3 Scale illustration

For illustration only, set $\alpha = \eta = 0.1$ and $\beta = 1/2$. The table reports the negative base-ten logarithm of the coefficient multiplying $\beta^{15} N^{17k+15}$, and the base-ten logarithm of the new sufficient prime-field threshold. These are formula evaluations, not experiments constructing the selected subset.

k	s	Old $-\log_{10}$ loss	New $-\log_{10}$ loss	New $\log_{10} N$ threshold
1	32	1.788×10^{11}	126.38	270.17
2	64	1.536×10^{21}	273.09	573.38
3	128	5.667×10^{40}	585.75	1218.10

The improvement is enormous locally, while the new field thresholds are still far beyond a practical exhaustive construction. The theorem is a quantitative structural statement, not a claim of an efficient large-instance algorithm.

II.171 Vector targets and exact torsion dependence

Source 40, delivered Section 8.

II.171.1 Dimension-free simultaneous selection

Corollary II.171.1 (Vector-valued maps). *Theorems F.45.1 and II.169.2 hold, with unchanged constants, for $\phi : B \rightarrow \mathbb{F}_q^m$ for every $m \geq 1$, when respected means that the vector defect is zero. Consequently the same bounds hold for simultaneously respecting any finite family of scalar maps, provided the hypothesis counts arrangements simultaneously respected by all of them.*

Proof. Choose $t \in \mathbb{F}_q^m$ uniformly and replace $t\phi(v)$ in (II.168.8) by the dot product $t \cdot \phi(v)$. Orthogonality now gives $\mathbb{E}_t \chi(t \cdot z) = \mathbf{1}_{z=0}$ for $z \in \mathbb{F}_q^m$ by taking the product of the coordinate identities.

All remaining steps are unchanged. A nonzero scalar $a \in \mathbb{F}_q$ cannot annihilate a nonzero vector defect. The only dimension dependence is the number of field coordinates in the random seed, not any counting constant. $\square$

The hypothesis is essential: knowing that each map separately respects many arrangements need not give a comparable intersection of their good arrangement families. The corollary does not manufacture this missing intersection.

II.171.2 General finite abelian targets

Let A be a finite abelian group, written additively, and $\phi : B \rightarrow A$. Define the defect by the same integer signed sum. Choose a character ψ of A uniformly from its character group $\hat{A}$ and choose $u \in \mathbb{F}_q$ uniformly. Use the probability

$$K_L(\psi(\phi(v))\chi(uF(v))). \quad (\text{II.171.1})$$

Character orthogonality on A says $\mathbb{E}_{\psi \in \hat{A}} \psi(a) = \mathbf{1}_{a=0}$; it follows by writing a finite abelian group as a product of cyclic groups and applying the finite geometric sum on each factor.

Proposition II.171.2 (Defect-order survival formula). *Assume $p > 2(L-1)$. For an $(L-1)$ -regular arrangement whose defect in A has order ν , its survival probability under (II.171.1) is exactly*

$$b_{s,L} R_{s,L}^{(\nu)}, \quad R_{s,L}^{(\nu)} = 1 + 2 \sum_{m=1}^{\lfloor (L-1)/\nu \rfloor} \left(1 - \frac{m\nu}{L}\right)^s. \quad (\text{II.171.2})$$

For a zero defect, take $\nu = 1$, so the probability is $a_{s,L}$. For a nonzero defect of order $\nu > L-1$, it is $b_{s,L}$.

Proof. The feature orthogonality still restricts coefficient vectors to $\eta = a\eta_0$. Target-character orthogonality then requires $a\Delta_\phi = 0$ in A , or equivalently $\nu \mid a$. Summing c_a^s over those multiples gives (II.171.2). Regularity gives distinct vertices, just as before. $\square$

Corollary II.171.3 (Targets without small torsion). *The selection theorems hold for $\phi : B \rightarrow A$ with unchanged constants provided A has no nonzero element killed by any integer $1 \leq a \leq L-1$. In particular, it suffices that every prime divisor of $|A|$ exceed $L-1$.*

Proof. Every nonzero defect then has order greater than $L-1$, so Proposition II.171.2 supplies exactly the bad probability $b_{s,L}$. Good arrangements, including repeated ones, retain the same lower bound by nonnegative coefficients. The score argument is unchanged. The sufficient condition on prime divisors implies that multiplication by each integer $1, \dots, L-1$ is an automorphism of A . $\square$

Remark II.171.4 (A concrete small-torsion obstruction). For fixed s and $\nu \geq 2$, Riemann sums in (II.171.2) show that

$$\frac{R_{s,L}^{(\nu)}}{R_s(L)} \longrightarrow \frac{1}{\nu} \quad (L \rightarrow \infty).$$

Thus the good-to-bad survival ratio against a defect of fixed order ν tends only to ν , rather than increasing without bound. One cannot extend the one-stage high-bandwidth conclusion to arbitrary small-torsion targets by reusing the same proof. This is an exact obstruction for this kernel calculation, not a nonexistence theorem for all selection methods.

II.172 A sharp exponent obstruction for positive Fourier kernels

Source 40, delivered Section 9.

This section explains in which precise sense the powers s and $s - 1$ are natural. It does not claim a universal extremal theorem about all subsets or all probabilistic constructions.

Definition II.172.1 (The homogeneous positive-Fourier class). Let

$$W(e^{i\theta}) = \sum_{a=-H}^H c_a e^{ia\theta}$$

be a trigonometric polynomial with real coefficients $c_{-a} = c_a \geq 0$, with $c_0 > 0$, and $0 \leq W \leq 1$ on the unit circle. For a single signed relation using every one of s distinct vertices, define the intended-relation good and bad probabilities by

$$a(W) = \sum_{a=-H}^H c_a^s, \quad b(W) = c_0^s, \quad R(W) = a(W)/b(W). \quad (\text{II.172.1})$$

These are the exact probabilities on regular configurations in the same orthogonality calculation as above. The same kernel is used at every vertex; this is the meaning of homogeneous here.

Theorem II.172.2 (Gain-mass obstruction). *For every kernel in this class and every integer $s \geq 2$,*

$$R(W) \leq \frac{1}{c_0}, \quad a(W) \leq b(W)^{(s-1)/s}, \quad a(W) \leq R(W)^{-(s-1)}. \quad (\text{II.172.2})$$

Consequently a required gain $R(W) \geq Q \geq 1$ forces $a(W) \leq Q^{-(s-1)}$. The same conclusion holds for any finite sequence of independent stages from this class: if their combined gain is at least Q , their combined intended-relation good survival probability is at most $Q^{-(s-1)}$.

Proof. Fourier inversion and $W \geq 0$ imply

$$0 \leq c_a \leq \frac{1}{2\pi} \int_0^{2\pi} W(e^{i\theta}) d\theta = c_0.$$

Also $\sum_a c_a = W(1) \leq 1$. Hence

$$\sum_a c_a^s \leq c_0^{s-1} \sum_a c_a \leq c_0^{s-1}.$$

Dividing by c_0^s gives $R(W) \leq 1/c_0$, while the same inequality is $a(W) \leq b(W)^{(s-1)/s}$. Since $c_0 \leq 1/R(W)$, it also gives the third inequality. If $R(W) \geq Q$, this is at most $Q^{-(s-1)}$.

For independent stages $W_1, \dots, W_r$, the regular good and bad probabilities multiply, so the gains multiply as well. Apply the third inequality to each stage and multiply:

$$\prod_i a(W_i) \leq \left(\prod_i R(W_i) \right)^{-(s-1)}.$$

□

[write] **7 October 2026; sources 40–46.** This barrier concerns the scalar first-moment certificate within the homogeneous positive-Fourier class. It advances source 04's and source 17's retention questions (item II.7) and sits beside source 17's Proposition II.62.2, which settles the scalar certificate in another class; it is not a lower bound for every selection method, as source 40 says.

Proposition II.172.3 (Fejér kernels attain the exponent). *For fixed $s \geq 2$ and any $Q \geq 1$, there is a Fejér kernel with $R_s(L) \geq Q$ and*

$$a_{s,L} \geq \frac{2^s}{(s+1)(s+3)^{s-1}} Q^{-(s-1)}. \quad (\text{II.172.3})$$

Proof. Take $L = \lceil (s+1)Q/2 \rceil \geq 2$. Then $R_s(L) \geq 2L/(s+1) \geq Q$, and $L \leq (s+3)Q/2$ because $Q \geq 1$. Thus

$$a_{s,L} \geq \frac{2}{s+1} L^{1-s} \geq \frac{2^s}{(s+1)(s+3)^{s-1}} Q^{1-s}.$$

□

II.172.1 Exactly what is optimal

Theorems II.172.2 and Proposition II.172.3 determine the power $Q^{-(s-1)}$ in the gain–good–mass trade-off, up to constants depending on s , within the stated class. A first-moment purity certificate using only the lower bound αT_0 for the initial good mass and an upper bound T_0 for the total mass requires gain of order $1/(\alpha\eta)$. In that certificate, the quantity $\alpha a(W)$ therefore has natural scale $\alpha^s \eta^{s-1}$, and the Fejér choice attains those powers.

This is not a lower bound on the actual number of good arrangements in the best selected subset. It does not exclude improvements from negative Fourier coefficients with a different positivity argument, vertex-dependent kernels, adaptive correlated stages, additional geometric information, higher-rank relation families, or a method that avoids the coarse T_0 comparison altogether. It also does not optimize the s -dependent constant or the field-size threshold.

II.173 Source 44’s chapter: the grid core of Lemma 15.4 [write]

Added in the eighth write, 7 October 2026.

The next nine sections are source 44’s Sections 2–10, on the count of degenerate arrangements in Gowers’s Lemma 15.4. *Logically the chapter follows the degeneracy material of sources 18, 22, 25 and 27* (Sections II.65 and II.81): source 18 found the leading term, source 25 the codimension-one classification and source 27 the core directions for $p > 2k + 2$, the second coefficient for $d = 2$ and the second coefficient for general d up to one unknown statistic $b_{2,d,p}$; it is printed here because no number printed earlier may move. Source 44 calls sources 25 and 27 “Baseline A” and “Baseline B” (note in Section F.49). Its introduction is Section F.49; its verification, questions and reproduction appendix are in Part V.

What is proved. In every characteristic $p \geq 5$ the core directions have width two, so the core count $k(2^{2d} - 1) + 2\binom{k}{2}$ holds without source 27’s condition $p > 2k + 2$ (Lemma II.175.1, Corollary II.175.2); the union of all core events is counted exactly, by $\prod_{i < k} (1 - (2^{2d} - 1 + 2i)/q^2)$ for $p \geq 5$ and $\prod_{i < k} (1 - p^{2d-1+i}/q^2)$ for $p = 2, 3$ (Theorems II.176.2, II.176.3), with an explicit non-core remainder of order q^{-3} (Theorem II.177.2); the second Whitney coefficient $b_{2,d}$ of the cross-section arrangement is characteristic-independent for $p > \max(2d, 16)$ (Theorem II.178.1) and is computed for every d by a weighted double count over subsets of the nine-point grid $\{-1, 0, 1\}^2$, in $O(d^3)$ operations (Theorem II.179.1), compressed by a finite rational certificate to eight lattice-walk energies (Theorem II.180.1). At $d = 8$, $b_{2,8} = 3,280,234,631,541$, so the second coefficient of the degenerate count is $\binom{k}{2} - 2,519,909k - 3,280,234,631,541 (2^{16} - 1 - H_8 = -2,519,909, H_8 = 2,585,444,$

source 18's C_8 ; checked at the write). The deficit of $b_{2,d}$ from $\binom{H_d}{2}$ has order $49^d/d$ (Theorem II.181.2). The intake reproduced $b_{2,d} = 3, 889, 97,081$ for $d = 2, 3, 4$ by brute force over hyperplane pairs (3 is source 27's $b_{2,2}$); $b_{2,8}$ rests on source 44's certificate, whose run was repeated at the write with identical output. This answers source 27's two questions (item II.38). The results are fixed- k asymptotics; core avoidance is not full nondegeneracy.

Relation to the formal statements (checked at HEAD 6d4942222). `lemma_15_4` (Sections14_15.lean, line 402) has the exact companion `lemma_15_4_holds` (Proofs15DegenerateArrangements.lean, line 1131), bound $3^{2d2^k} k N^{(2d+1)k+2d-2}$; source 44 follows its degeneracy definition (`GeneralArrangement.IsDegenerate`, line 354; `degenerateGeneralArrangementCount`, line 364), whose file has the blob it records. Nothing in this chapter is formalized.

II.174 Arrangements, coefficient directions, and the regular region

Source 44, delivered Section 2.

Fix $d \geq 2$, $k \geq 1$, and a finite field $\mathbb{F}_q$ of characteristic p . Write

$$m = 2d, \quad v = m2^k, \quad D = (m+1)k + m - 1, \quad w = (\underbrace{1, \dots, 1}_d, \underbrace{-1, \dots, -1}_d). \quad (\text{II.174.1})$$

The signs are interpreted in the field, even when $p = 2$. Since $w \neq 0$,

$$W = \{r \in \mathbb{F}_q^m : w \cdot r = 0\}$$

has dimension $m - 1$. A labelled arrangement consists of $h \in \mathbb{F}_q^k$, $x = (x_{ji}) \in \mathbb{F}_q^{m \times k}$, and $r \in W$, with vertices

$$\rho(\epsilon, j) = (x_{j1} + \epsilon_1 h_1, \dots, x_{jk} + \epsilon_k h_k, r_j), \quad \epsilon \in \{0, 1\}^k, \quad 1 \leq j \leq m. \quad (\text{II.174.2})$$

There are exactly q^D parameter tuples. No distinctness of vertices is assumed in this count.

Let $\mathcal{A}_p$ be the image of $\{-1, 0, 1\}$ in $\mathbb{F}_p$. For an array $\eta \in \mathcal{A}_p^{\{0,1\}^k \times [m]}$ and $A \subseteq [k+1]$, define

$$M_A(\eta) = \sum_{j, \epsilon} \eta(\epsilon, j) \prod_{i \in A} \rho(\epsilon, j)_i.$$

The empty set is included. Put

$$e(\epsilon, j) = (-1)^{\epsilon_1 + \dots + \epsilon_k} w_j.$$

The alternating cube identity, together with $\sum_j w_j = 0$ and $w \cdot r = 0$, proves $M_A(e) = 0$ for every A .

Definition II.174.1. An arrangement is *degenerate* if there exists such an η outside $\mathbb{F}_q e$ for which every $M_A(\eta)$ vanishes. Write $\Delta_{k,d}(q) = Z_{k,d}(q)/q^D$. Two admissible non-intrinsic arrays are in the same *direction* if $\eta' = a\eta + b e$ for some $a \neq 0$.

Equivalent arrays have exactly the same moment-zero event. Dividing pointwise by e turns the equivalence into affine equivalence of profiles $t' = at + b\mathbf{1}$. If nonconstant profiles take prime-field values, their relating scalars a, b lie in the prime field: two unequal coordinates determine a , and either coordinate then determines b .

An array is *column-parity* if $\eta(\epsilon, j) = (-1)^{|\epsilon|} a_j$ for every j . When all h_i are nonzero, its moment equations are precisely

$$\sum_j a_j = 0, \quad a \cdot r = 0. \quad (\text{II.174.3})$$

Here $a \in \mathcal{A}_p^m$, and a must not be proportional to w . Let $\mathcal{H}_{d,p}$ be the distinct hyperplanes in W so obtained, and set

$$W^\circ = W \setminus \bigcup_{H \in \mathcal{H}_{d,p}} H, \quad g_{d,p}(q) = \frac{|W^\circ|}{q^{m-1}}, \quad H_{d,p} = |\mathcal{H}_{d,p}|. \quad (\text{II.174.4})$$

Lemma II.174.2. *Every arrangement with some $h_i = 0$ is degenerate. If $r \in W^\circ$, its coordinates are pairwise distinct. The complement of the region $h_i \neq 0$ for all i and $r \in W^\circ$ is entirely degenerate.*

Proof. A zero side gives two equal labelled vertices. Coefficients $1, -1$ at these labels and zero elsewhere annihilate all moments, but cannot be a nonzero multiple of the full-support array e . This works in characteristic two as well. For distinct $u, t \in [m]$, the vector $a = \mathbf{e}_u - \mathbf{e}_t$ is alphabet-valued, has sum zero, and is not proportional to w , since $m \geq 4$. Thus $r_u = r_t$ is one of the forbidden hyperplanes. The final assertion follows from (II.174.3). $\square$

II.175 A width theorem with no dimension-dependent characteristic cutoff

Source 44, delivered Section 3.

The core directions from [200] can be written as

$$\eta(\epsilon, j) = e(\epsilon, j) \left(t_j + \sum_{i=1}^k b_i \epsilon_i \right), \quad b \neq 0, \quad (\text{II.175.1})$$

with all displayed profile values in $\mathcal{A}_p$. In particular, the slopes b_i are common to the columns. The elementary lemma below removes the earlier requirement $p > 2k + 2$.

Lemma II.175.1 (Characteristic-uniform width two). *Let $p \geq 5$. A nonconstant affine Boolean function $f(\epsilon) = t + \sum_i b_i \epsilon_i$ with values in $\mathcal{A}_p$ has at most two nonzero slopes. With one active coordinate, the two endpoint values are any two distinct alphabet elements. With two active coordinates, the slopes are ± 1 , and the offset is forced so that the three attained values are $-1, 0, 1$.*

Proof. The value at zero and every single-coordinate difference belong to $\mathbb{F}_p$, so all parameters are in that field. If A is a nonempty proper subset of $\mathbb{F}_p$ and $b \neq 0$, then $A + b \neq A$: otherwise translation by b would make A invariant under the whole additive group. Consequently $|A \cup (A + b)| \geq |A| + 1$. Adding active slopes successively grows the attained-value set by at least one until it fills the field. Three active slopes would therefore give at least four values, which is impossible when $p \geq 5$.

For two active slopes b, c , the four values are $t, t + b, t + c, t + b + c$. A collision, with both slopes nonzero, forces $b = c$ or $b = -c$. The resulting three values form an arithmetic progression. If they are exactly the ternary alphabet, their middle value is zero: their sum is three times that middle value, and $3 \neq 0$. The step is therefore 1 or -1 . The one-coordinate case requires only that its two endpoint values be distinct alphabet elements. $\square$

[write] **7 October 2026; sources 40–46.** This removes the condition $p > 2k + 2$ of source 27's core classification (Theorem II.97.6), answering source 27's question on intermediate characteristics $5 \leq p \leq 2k + 2$.

Corollary II.175.2. *For every $p \geq 5$, every $d \geq 2$, and every $k \geq 1$, the number of core directions is*

$$S_{k,d} = k(2^m - 1) + 2 \binom{k}{2}. \quad (\text{II.175.2})$$

Thus no new core families occur for $5 \leq p \leq 2k + 2$.

Proof. For a fixed single-coordinate slope direction, normalize a unit slope to 1. Each offset can be chosen in $\{-1, 0\}$. The two constant offset patterns are equivalent under common translation; all other binary patterns remain distinct. This gives $2^m - 1$ directions. A slope of absolute value two forces constant offsets and yields the already counted constant-offset direction after affine rescaling.

For two active coordinates, the normalized slope ratio is 1 or -1 . The width lemma forces a common offset, which disappears modulo constant translation. These give two directions for each coordinate pair. Distinct supports or slope ratios cannot be related by a single affine change of profile values. $\square$

This answers the intermediate-characteristic question in [200] without assuming that an integer lift of a long Boolean sum has small absolute value. The translation argument works directly in the prime field. Characteristics two and three are genuinely different: the alphabet is then the entire prime field and allows arbitrarily many active slopes.

II.176 Exact core-avoidance laws

Source 44, delivered Section 4.

Fix nonzero sides and a regular cross-section vector r . Define the paired linear map

$$L_r(t) = \left(\sum_j w_j t_j, \sum_j w_j r_j t_j \right), \quad Z_i = L_r((x_{ji}/h_i)_{j=1}^m) \in \mathbb{F}_q^2. \quad (\text{II.176.1})$$

Its two coefficient rows, w and $w \odot r$, are independent because r is not constant. Thus $Z_1, \dots, Z_k$ are independent uniform points of $\mathbb{F}_q^2$ when the bases x are uniform.

For a core profile (II.175.1), all but the two top moments vanish automatically. Expansion of those top moments gives the equivalent condition

$$\sum_i b_i Z_i = -L_r(t). \quad (\text{II.176.2})$$

Indeed each term ϵ_i contributes $1 + x_{ji}/h_i$ after division by the alternating top-cube factor. The extra common constant $\sum_i b_i$ vanishes under L_r , since $L_r(\mathbf{1}) = 0$.

II.176.1 Characteristics at least five: signed collisions

Put $L = 2^m - 1$ and define the binary offset set

$$\mathcal{C}_r = \{L_r(u) : u \in \{0, 1\}^m\} \subseteq \mathbb{F}_q^2. \quad (\text{II.176.3})$$

Lemma II.176.1. *For $p \geq 5$ and $r \in W^\circ$, the set $\mathcal{C}_r$ has exactly L elements, contains zero, and is invariant under negation.*

Proof. The all-zero and all-one vectors have the same image, namely zero. If two binary vectors u, u' have the same image, then $a = w \odot (u - u')$ is ternary, has sum zero, and satisfies $a \cdot r = 0$. Regularity forces $a = \lambda w$, hence $u - u' = \lambda \mathbf{1}$. Since a binary difference is 0, 1 or -1 , a nonzero constant difference occurs only for the all-one and all-zero pair. These are the only duplicate images. Finally $L_r(\mathbf{1} - u) = -L_r(u)$. $\square$

Theorem II.176.2 (Exact signed-collision product). *For every $p \geq 5$, nonzero h , and regular r , the probability that no core direction vanishes is*

$$F_{k,d,p}(q) = \prod_{i=0}^{k-1} \left(1 - \frac{2^m - 1 + 2i}{q^2} \right). \quad (\text{II.176.4})$$

Equivalently, the exact number of such base matrices is

$$q^{(m-2)k} \prod_{i=0}^{k-1} (q^2 - L - 2i). \quad (\text{II.176.5})$$

Proof. The single-coordinate core events are exactly $Z_i \in \mathcal{C}_r$ by (II.176.2) and the proof of Corollary II.175.2. The two-coordinate events are $Z_i = Z_j$ or $Z_i = -Z_j$. Outside $\mathcal{C}_r$, negation has no fixed points, because $0 \in \mathcal{C}_r$ and p is odd. There are $(q^2 - L)/2$ two-element sign orbits. Choose an ordered list of k distinct orbits and one of the two signs in each: the number of possible paired tuples is $\prod_{i < k} (q^2 - L - 2i)$. Each paired tuple has $q^{(m-2)k}$ preimages under the independent rank-two maps in (II.176.1). $\square$

If k exceeds the number of available orbits, the product is zero: a zero factor occurs before any negative factors. Thus the formula also covers the saturation range. If there are no regular cross-sections, the conditional statement is vacuous and the unconditional regular-region count below is zero.

II.176.2 Characteristics two and three: linear independence

Theorem II.176.3 (Exact prime-field rank product). *Let $p \in \{2, 3\}$ and $q = p^a$. For fixed nonzero sides and regular r ,*

$$F_{k,d,p}(q) = \prod_{i=0}^{k-1} \left(1 - \frac{p^{m-1+i}}{q^2} \right) \quad (\text{II.176.6})$$

is the exact probability that no core direction vanishes. Moreover,

$$g_{d,p}(q) = \prod_{j=0}^{m-3} \left(1 - \frac{p^j}{q} \right). \quad (\text{II.176.7})$$

Proof. The allowed offsets form all of $\mathbb{F}_p^m$. Let $C_r = L_r(\mathbb{F}_p^m)$, an $\mathbb{F}_p$ -subspace of $\mathbb{F}_q^2$. Its kernel consists exactly of the constant vectors. To see this, $L_r(t) = 0$ says that $c = w \odot t \in \mathbb{F}_p^m$ has sum zero and $c \cdot r = 0$. Since the alphabet is all of $\mathbb{F}_p$, regularity implies $c \in \mathbb{F}_p w$, or $t \in \mathbb{F}_p \mathbf{1}$. Therefore

$$\dim_{\mathbb{F}_p} C_r = m - 1.$$

A core relation exists precisely when a nonzero vector $b \in \mathbb{F}_p^k$ satisfies $\sum_i b_i Z_i \in C_r$. The images of the Z_i in the quotient $\mathbb{F}_q^2 / C_r$ must therefore be linearly independent. This quotient has dimension

$2a - m + 1$. After i independent images have been chosen, the preimage of their span contains p^{m-1+i} of the q^2 possible next summaries. Multiplying the successive avoidance probabilities proves (II.176.6).

For (II.176.7), put

$$E_0 = \{c \in \mathbb{F}_p^m : \sum_j c_j = 0\} / \mathbb{F}_p w, \quad \dim_{\mathbb{F}_p} E_0 = m - 2.$$

A cross-section $r \in W$ induces the $\mathbb{F}_p$ -linear map $E_0 \rightarrow \mathbb{F}_q$, $c \mapsto c \cdot r$. The map from W onto the space of all such linear maps is surjective and has kernel $\mathbb{F}_q \mathbf{1}$. Indeed the annihilator of all coordinate-sum-zero vectors is the constant line, and the dimensions over $\mathbb{F}_q$ then give surjectivity. Regularity is exactly injectivity of the induced map on E_0 . The probability that $m - 2$ uniform elements of $\mathbb{F}_q$ are $\mathbb{F}_p$ -linearly independent is the product in (II.176.7). $\square$

Corollary II.176.4 (Universal obstructions). *In characteristics two and three, every arrangement is degenerate if $a < m - 2$ or $k > 2a - m + 1$. In characteristic at least five, every arrangement is degenerate if $k > (q^2 - 2^m + 1)/2$.*

Proof. Nondegeneracy requires nonzero sides, a regular cross-section, and core avoidance. The required injections or distinct sign orbits cannot exist beyond the displayed dimensions. If no regular cross-section exists, the conclusion follows already from Lemma II.174.2. $\square$

These are sufficient conditions for universal degeneracy, or equivalently necessary conditions for the existence of a nondegenerate arrangement. Their converses are not asserted.

II.177 From exact core avoidance to the full degeneracy count

Source 44, delivered Section 5.

The product formula is not, by itself, a formula for full nondegeneracy. Other coefficient directions remain. This section retains the cubic remainder from the earlier analysis [200], proves the needed rank statement, and substitutes the exact product in place of the earlier pairwise estimate.

II.177.1 The exact inventory of coefficient directions

For even ℓ , define

$$T_\ell(p) = \sum_{a \equiv 0 \pmod{p}} [z^a] (z^{-1} + 1 + z)^\ell, \quad (\text{II.177.1})$$

$$B_\ell(p) = \sum_{j \in \mathbb{Z}} \binom{\ell}{\ell/2 + jp}. \quad (\text{II.177.2})$$

A binomial coefficient is zero outside its usual range. Let $Q_\ell(p)$ count nonconstant affine classes of alphabet-valued profiles with balanced signed coordinate sum, and let $U_\ell(p)$ count such classes without the balance requirement. Then

$$Q_\ell(p) = \frac{T_\ell(p) - 2B_\ell(p) + 1}{2}, \quad U_\ell(p) = \frac{3^\ell - 2^{\ell+1} + 1}{2} \quad (p \geq 5), \quad (\text{II.177.3})$$

$$Q_\ell(p) = \frac{p^{\ell-2} - 1}{p - 1}, \quad U_\ell(p) = \frac{p^{\ell-1} - 1}{p - 1} \quad (p = 2, 3). \quad (\text{II.177.4})$$

Here the balanced signs have equally many 1 and -1 entries.

For completeness, when $p \geq 5$ there are $T_\ell(p)$ balanced ternary profiles and $B_\ell(p)$ balanced binary profiles. There are three constant ternary profiles. Profiles using exactly two values give $(B_\ell(p) - 2)/2$ affine classes, because the underlying subset is identified with its complement. Profiles using all three values have only identity and negation as affine alphabet symmetries; they give $(T_\ell(p) - 3B_\ell(p) + 3)/2$ classes. Their sum is (II.177.3). Removing balance replaces these two profile counts by 3^ℓ and 2^ℓ . In characteristics two and three the alphabet is the whole field. Quotient the vector space of profiles first by the constant line and then projectivize; this proves (II.177.4).

Consequently the relevant inventory is

$$H_{d,p} = Q_m(p), \quad R_{k,d,p} = Q_v(p) - U_m(p), \quad (\text{II.177.5})$$

where $R_{k,d,p}$ is the number of non-column-parity directions satisfying the empty moment. Indeed column-parity profiles are arbitrary profiles on m columns and automatically have zero empty moment. The core inventory is

$$S_{k,d,p} = \begin{cases} k(2^m - 1) + k(k - 1), & p \geq 5, \\ p^{m-1}(p^k - 1)/(p - 1), & p = 2, 3. \end{cases} \quad (\text{II.177.6})$$

The second line follows by choosing a projective nonzero slope in $\mathbb{F}_p^k$ and an offset modulo the constant line. In particular $R_{k,d,p} - S_{k,d,p}$ is nonnegative: it counts the remaining directions. The inventory formulae themselves are inherited from [200]; our improvement is that the large-characteristic core formula now holds for every $p \geq 5$.

II.177.2 Why every remaining direction costs three powers of the field size

For $A \subseteq [k]$ write

$$c_{A,j} = \sum_{\epsilon} \eta(\epsilon, j) \prod_{i \in A} \epsilon_i, \quad h_A = \prod_{i \in A} h_i.$$

The block moment satisfies the exact Boolean expansion

$$F_{A,j} := \sum_{\epsilon} \eta(\epsilon, j) \prod_{i \in A} (x_{ji} + \epsilon_i h_i) = \sum_{B \subseteq A} h_B c_{B,j} \prod_{i \in A \setminus B} x_{ji}. \quad (\text{II.177.7})$$

The two full moments associated with A are $\sum_j F_{A,j}$ and $\sum_j r_j F_{A,j}$. Boolean inversion, valid in every characteristic, shows that column-parity is equivalent to the vanishing of all vectors c_A with $A \subsetneq [k]$.

Lemma II.177.1 (Non-core remainder). *For a fixed non-column-parity direction which is not core, conditional on nonzero sides, the probability that its moments vanish and the cross-section is regular is at most q^{-3} .*

Proof. Let $s < k$ be the smallest degree for which some $c_A \neq 0$, $|A| = s$. The degree- s moments force

$$\sum_j c_{A,j} = 0, \quad c_A \cdot r = 0. \quad (\text{II.177.8})$$

If the first condition fails there are no solutions. For each $i \notin A$, (II.177.7) at $A \cup \{i\}$ is affine in the coordinate block $(x_{1i}, \dots, x_{mi})$, with coefficient vector $h_A c_A$. Its paired last-coordinate moment has

coefficient vector $h_A(r \odot c_A)$. A nonzero vector of sum zero has at least two nonzero coordinates. Since regular r has pairwise distinct entries, these two rows have rank two.

If some minimal-degree c_A is not proportional to w , its cross-section constraint in (II.177.8) has codimension one in W . For each admissible regular r , the two base equations cost a further factor q^{-2} . Their joint probability is at most q^{-3} .

Suppose instead that every minimal-degree vector is proportional to w . If $s \leq k - 2$, choose $i, \ell \notin A$ distinct. The paired moments at $A \cup \{i\}$ and $A \cup \{\ell\}$ give four independent equations: the first pair has its rank-two rows in the i -coordinate block and the second in the ℓ -coordinate block. Their other terms depend only on the A -coordinates and the sides. Fixing those coordinates therefore gives probability at most q^{-4} , which is sufficient.

The remaining case has $s = k - 1$ and every proper coefficient vector proportional to w . Boolean inversion then says precisely $\eta/e = t_j + \sum_i b_i \epsilon_i$ with common slopes and $b \neq 0$. This is a core direction, contrary to the hypothesis. These cases exhaust the possibilities. $\square$

The assertion about the other terms is important. In (II.177.7), at $A \cup \{i\}$, every term involving x_{ji} has coefficient c_B with $B \subseteq A$. Minimality leaves only $B = A$. All other surviving terms use coordinates from A , not from another block outside A . Thus the rank-four step does not assume independence of unexamined right sides.

Let $F_{k,d,p}(q)$ be the product in Theorem II.176.2 or II.176.3, as appropriate, and put

$$\Phi_{k,d,p}(q) = (1 - q^{-1})^k g_{d,p}(q) F_{k,d,p}(q). \quad (\text{II.177.9})$$

When there are no regular cross-sections set $\Phi = 0$; this agrees with the product expression because $g = 0$.

Theorem II.177.2 (Exact principal part and explicit cubic remainder). *In every characteristic,*

$$0 \leq \Delta_{k,d}(q) - (1 - \Phi_{k,d,p}(q)) \leq (1 - q^{-1})^k \frac{R_{k,d,p} - S_{k,d,p}}{q^3}. \quad (\text{II.177.10})$$

The quantity Φ is the exact probability of simultaneously avoiding zero sides, all cross-section hyperplanes, and all core relations.

Proof. Every excluded zero side or cross-section is degenerate. On the remaining region, all core events have the exact avoidance probability already proved, independently of the particular regular r and nonzero h . This proves the interpretation of Φ and the lower inequality. Every additional degeneracy has a non-core witness. Sum Lemma II.177.1 over the $R - S$ directions, then multiply by the probability of nonzero sides. This proves the upper inequality. $\square$

For example, the earlier S/q^2 upper bound for the core union discards its exact intersections. The replacement is $1 - \prod_{i < k} (1 - (2^m - 1 + 2i)/q^2)$ in characteristics at least five. The improvement is an exact finite identity, not a claim that the non-core term is uniformly small when k grows. A convenient inherited coarse consequence remains

$$\Delta_{k,d}(q) \leq \frac{k + H_{d,p}}{q} + \frac{S_{k,d,p}}{q^2} + \frac{R_{k,d,p} - S_{k,d,p}}{q^3}. \quad (\text{II.177.11})$$

The exact expression (II.177.10) is preferable whenever the cross-section complement can itself be evaluated.

II.178 The second Whitney coefficient and a uniform characteristic cutoff

Source 44, delivered Section 6.

Work first over $\mathbb{Q}$. The cross-section arrangement is defined by the integer-balanced ternary vectors $a \in \{-1, 0, 1\}^m$, $\sum_j a_j = 0$, modulo the intrinsic line $\mathbb{Q}w$. Let $\mathcal{L}_2$ be its distinct codimension-two flats in W , and let ν_F be the number of hyperplanes containing F . Define

$$b_{2,d} = \sum_{F \in \mathcal{L}_2} (\nu_F - 1). \quad (\text{II.178.1})$$

This is the second Whitney coefficient of the characteristic polynomial; we will use only the finite inclusion–exclusion interpretation, not any structural theorem about arrangements.

Put

$$T_n = [z^0](z^{-1} + 1 + z)^n, \quad H_d = \frac{T_m - 2\binom{m}{d} + 1}{2}. \quad (\text{II.178.2})$$

The profile count (II.177.3) gives H_d in characteristic zero or greater than m . It does not alone establish that pair intersections have stabilized. The following determinant bound does.

Theorem II.178.1 (Rank-two stability beyond characteristic sixteen). *If $p > \max(m, 16)$, reduction from the rational arrangement to $\mathcal{H}_{d,p}$ preserves the hyperplanes, every rank-two flat, and every incidence between them. In particular $H_{d,p} = H_d$ and $b_{2,d,p} = b_{2,d}$.*

Proof. Eliminate r_m using $w_m = -1$ and $w \cdot r = 0$. The unnormalized restricted integer normal is

$$n(a) = (a_i + a_m w_i)_{1 \leq i < m}. \quad (\text{II.178.3})$$

For $t \leq 3$ normals $n(a^{(1)}), \dots, n(a^{(t)})$, each t by t minor is, up to sign, the determinant of the $(t+1)$ by $(t+1)$ matrix whose rows are $a^{(1)}, \dots, a^{(t)}, w$ restricted to the chosen t coordinates and coordinate m . This follows by adding suitable multiples of the last row to the preceding rows and expanding at the last column. All entries of this augmented matrix are in $\{-1, 0, 1\}$. Hadamard's inequality therefore bounds the absolute determinant by

$$(t+1)^{(t+1)/2} \leq 16.$$

A nonzero minor cannot become zero modulo $p > 16$; a zero integer minor remains zero. All ranks of collections of at most three normals are preserved. Nonzero normals and proportional pairs are consequently preserved. Membership of a third normal in the span of two independent ones is preserved as well. These are exactly the data determining the rank-two flats and their incident hyperplanes.

Finally $p > m$ means that every modularly balanced ternary vector is integer-balanced, since its integer coordinate sum has absolute value at most m . Thus no new normals are introduced by a modular balance condition. This proves the whole assertion. $\square$

The bound 16 is independent of d and k because only three restricted normals are needed. It does not claim that the entire intersection lattice stabilizes at that prime. Higher-rank minors can have larger determinants.

Lemma II.178.2 (A uniform two-term complement expansion). *For $p > \max(m, 16)$,*

$$g_{d,p}(q) = 1 - \frac{H_d}{q} + \frac{b_{2,d}}{q^2} + O_d(q^{-3}), \quad (\text{II.178.4})$$

uniformly over the indicated finite fields. One valid absolute remainder constant is

$$C_d = (1 + H_d + b_{2,d}) \binom{H_d}{3}. \quad (\text{II.178.5})$$

Proof. At a point outside all intersections of three linearly independent hyperplane normals, the incident normals have rank at most two. If there are s incident hyperplanes of rank two, all of them contain one unique rank-two flat and its correction is $s - 1$. Thus, at every such point,

$$\mathbf{1}_{\cup H} = \sum_{H \in \mathcal{H}_{d,p}} \mathbf{1}_H - \sum_{F \in \mathcal{L}_2} (\nu_F - 1) \mathbf{1}_F.$$

The identity is also true at rank zero or one. The exceptional points belong to a union of at most $\binom{H_d}{3}$ codimension-three subspaces, of total relative size at most $\binom{H_d}{3} q^{-3}$. The discrepancy of the displayed indicator formula has absolute value at most $1 + H_d + b_{2,d}$. Averaging gives the asserted error bound. Every hyperplane has relative size q^{-1} and every rank-two flat relative size q^{-2} . Theorem II.178.1 makes their coefficients independent of the characteristic. $\square$

Corollary II.178.3 (The coefficient to be computed). *For fixed k, d , uniformly for $p > \max(2d, 16)$,*

$$\Delta_{k,d}(q) = \frac{k + H_d}{q} \quad (\text{II.178.6})$$

$$+ \frac{\binom{k}{2} + k(2^m - 1 - H_d) - b_{2,d}}{q^2} + O_{k,d}(q^{-3}). \quad (\text{II.178.7})$$

Proof. For fixed k, d , the exact core product is $1 - S_{k,d}/q^2 + O_{k,d}(q^{-4})$. Substitute this, the binomial expansion of $(1 - q^{-1})^k$, and (II.178.4) into (II.177.9). Theorem II.177.2 gives an $O_{k,d}(q^{-3})$ remainder uniformly in the characteristic: the direction inventory is bounded by 3^v , independently of p . Since $S_{k,d} - \binom{k}{2} - kH_d = \binom{k}{2} + k(2^m - 1 - H_d)$, the coefficient is the one displayed. $\square$

The dependence on $b_{2,d,p}$ in this corollary was already identified in [200]. The new content so far is its uniform characteristic range and the exact principal product. The next section evaluates the remaining invariant in a number of operations polynomial in d .

II.179 A fixed nine-point grid computes every second coefficient

Source 44, delivered Section 7.

II.179.1 Affine profiles and rank-two flats

For an integer-balanced ternary normal a , put $z = w \odot a \in \{-1, 0, 1\}^m$. Then $w \cdot z = 0$. Two nonconstant profiles define the same hyperplane exactly when

$$z' = \lambda z + \mu \mathbf{1}, \quad \lambda \neq 0. \quad (\text{II.179.1})$$

A profile using exactly two distinct values has six ternary representatives in its affine class. A profile using all three values has two, because the only affine symmetries of $\{-1, 0, 1\}$ over $\mathbb{Q}$ are identity and negation. These representative multiplicities will be retained, rather than replacing affine equivalence by sign equivalence.

Take an ordered pair of nonconstant balanced profiles z, z' which are independent modulo the constant line. The coordinate pairs form a set

$$S = \{(z_j, z'_j) : 1 \leq j \leq m\} \subseteq \mathcal{G} := \{-1, 0, 1\}^2. \quad (\text{II.179.2})$$

Independence modulo constants is equivalent to S affinely spanning $\mathbb{Q}^2$. Let $\mathcal{S}$ denote all subsets of $\mathcal{G}$ with this property. There are only $2^9 = 512$ possible subsets before imposing it. For $S \in \mathcal{S}$ define

$$c_i(S) = \begin{cases} 6, & \text{the } i\text{th projection has two values,} \\ 2, & \text{it has three values,} \end{cases} \quad i = 1, 2. \quad (\text{II.179.3})$$

Both projections have at least two values, since S spans the plane.

Let $\nu(S)$ be the number of nonzero projective directions $(b : c)$ for which some affine function $a + bX + cY$ maps S into $\{-1, 0, 1\}$. Every such function produces a balanced ternary profile $a\mathbf{1} + bz + cz'$, because the three constituent profiles have signed sum zero. Its class gives a hyperplane containing the codimension-two flat defined by the original pair. Conversely every hyperplane through that flat arises this way. Therefore $\nu(S)$ is exactly its multiplicity.

There is a small exact algorithm for $\nu(S)$. Choose any three noncollinear points of S . Assign to them the 27 ordered triples of values in $\{-1, 0, 1\}^3$, solve for the unique affine function, and retain it if its values at the other points are in $\{-1, 0, 1\}$. Discard constant functions and identify proportional nonzero gradient vectors. This is an exhaustive rational computation, since an affine function is determined by three noncollinear values.

II.179.2 Balanced coordinate sequences and lattice energies

For any $T \subseteq \mathcal{G}$, let

$$P_T(X, Y) = \sum_{(u,v) \in T} X^u Y^v, \quad E_d(T) = \text{CT} \left(P_T(X, Y)^d P_T(X^{-1}, Y^{-1})^d \right). \quad (\text{II.179.4})$$

This counts ordered pairs of length- d sequences from T with equal vector sum. Equivalently, if $W_d^T(s)$ counts ordered d -step walks with steps in T and endpoint $s \in \mathbb{Z}^2$, then

$$W_0^T(s) = \mathbf{1}_{s=0}, \quad W_{j+1}^T(s) = \sum_{t \in T} W_j^T(s - t), \quad E_d(T) = \sum_s W_d^T(s)^2. \quad (\text{II.179.5})$$

For $d \geq 1$, $E_d(\emptyset) = 0$. Define the exact-support count

$$N_d(S) = \sum_{T \subseteq S} (-1)^{|S| - |T|} E_d(T). \quad (\text{II.179.6})$$

Inclusion-exclusion shows that $N_d(S)$ counts the balanced coordinate sequences whose combined support is exactly S . In particular it is a nonnegative integer, despite the alternating formula.

Theorem II.179.1 (Fixed-grid formula). *For every $d \geq 2$,*

$$b_{2,d} = \sum_{S \in \mathcal{S}} \frac{N_d(S)}{\nu(S)c_1(S)c_2(S)}. \quad (\text{II.179.7})$$

More precisely, if $f_{\nu,d}$ is the number of rank-two flats incident with exactly ν hyperplanes, then

$$f_{\nu,d} = \sum_{\substack{S \in \mathcal{S} \\ \nu(S)=\nu}} \frac{N_d(S)}{\nu(\nu-1)c_1(S)c_2(S)}. \quad (\text{II.179.8})$$

Both formulae can be evaluated through d in $O(d^3)$ exact arithmetic operations and $O(d^2)$ memory, with constants independent of d .

Proof. Fix a rank-two flat F of multiplicity ν . It has $\nu(\nu-1)$ ordered pairs of distinct incident hyperplanes. For an ordered pair, choose every ternary representative of each of its two profiles. There are $c(z)c(z')$ raw ordered pairs. The support S of each raw pair spans the plane; its multiplicity is ν , and its two projection representative counts are $c(z), c(z')$.

Give this raw pair weight $1/(\nu c(z)c(z'))$. All representatives of one ordered hyperplane pair contribute $1/\nu$. All ordered hyperplane pairs through F therefore contribute $\nu-1$. On the other hand, raw balanced pairs with a specified support S are counted by $N_d(S)$: the first d coordinate pairs and the last d coordinate pairs must have equal sum. Summing the weights by support proves (II.179.7). Giving each raw pair instead weight $1/(\nu(\nu-1)c(z)c(z'))$ proves (II.179.8).

For the operation count, at step j every walk endpoint lies in $[-j, j]^2$. The recurrence (II.179.5) uses at most nine additions per endpoint, so one step costs $O(j^2)$ operations for each of the fixed 512 supports. Summing for $j \leq d$ costs $O(d^3)$ operations, and only the current arrays, of total size $O(d^2)$, need be stored. Boolean Möbius inversion on the nine-bit subset lattice costs a fixed number of operations for each j . All rational weights and template multiplicities are independent of d . The final numbers are integers by the already proved counting interpretation. $\square$

The bound is an arithmetic-operation bound, not a unit-cost assertion for arbitrarily large integers. Walk counts have $O(d)$ bits, since there are at most 9^d walks; the same linear bit bound holds for the energies. Using ordinary schoolbook integer arithmetic still gives a polynomial bit-complexity bound. The algorithm replaces enumeration of exponentially many hyperplanes or their quadratically many pairs by a fixed collection of two-dimensional walk recurrences.

The formula also gives an independently checkable consistency identity:

$$\sum_{\nu} \binom{\nu}{2} f_{\nu,d} = \binom{H_d}{2}, \quad b_{2,d} = \sum_{\nu} (\nu-1) f_{\nu,d}. \quad (\text{II.179.9})$$

The first identity says that every unordered pair of hyperplanes meets in one rank-two flat. Neither identity is used to guess the answer.

II.180 Eight lattice energies and the explicit eight-arrangement coefficient

Source 44, delivered Section 8.

II.180.1 A finite-certified compression of the general formula

Define the pair-overlap deficit

$$B_d = \binom{H_d}{2} - b_{2,d}. \quad (\text{II.180.1})$$

This notation is distinct from the modular binary count $B_\ell(p)$, which always has two arguments and was used only in the direction inventory. By (II.179.9),

$$B_d = \sum_{F \in \mathcal{L}_2} \binom{\nu_F - 1}{2} = \sum_{S \in \mathcal{S}} \frac{\nu(S) - 2}{2\nu(S)c_1(S)c_2(S)} N_d(S). \quad (\text{II.180.2})$$

The first equality makes positivity transparent. The second follows by the same weighted raw-pair count as Theorem II.179.1.

For a Laurent polynomial P define

$$\mathcal{E}_d(P) = \text{CT}(P(X, Y)^d P(X^{-1}, Y^{-1})^d).$$

Consider the following eight polynomials, each the indicator polynomial of a finite planar step set:

$$P_7 = 1 + X + X^{-1} + Y + Y^{-1} + XY^{-1} + X^{-1}Y, \quad (\text{II.180.3})$$

$$P_6 = 1 + X + Y + X^2 + XY + Y^2, \quad (\text{II.180.4})$$

$$P_{5a} = 1 + X + Y + XY + Y^2, \quad (\text{II.180.5})$$

$$P_{5b} = 1 + X + X^{-1} + Y + Y^{-1}, \quad (\text{II.180.6})$$

$$P_4 = 1 + X + Y + Y^2, \quad (\text{II.180.7})$$

$$P_3 = 1 + X + Y, \quad P_2 = 1 + X, \quad P_1 = 1. \quad (\text{II.180.8})$$

The subscripts record the number of points, not the arrangement dimension. In particular $\mathcal{E}_d(P_2) = \binom{2d}{d}$ and $\mathcal{E}_d(P_1) = 1$.

Theorem II.180.1 (Eight-shape compression; finite rational certificate). *For every $d \geq 1$, the fixed-grid expression gives*

$$\begin{aligned} B_d = & \frac{1}{12}\mathcal{E}_d(P_7) + \frac{1}{6}\mathcal{E}_d(P_6) - \frac{1}{2}\mathcal{E}_d(P_{5a}) + \frac{1}{8}\mathcal{E}_d(P_{5b}) \\ & - \frac{1}{2}\mathcal{E}_d(P_4) + \frac{7}{6}\mathcal{E}_d(P_3) - \binom{2d}{d} + \frac{19}{24}. \end{aligned} \quad (\text{II.180.9})$$

For $d = 1$ both H_d and $b_{2,d}$ are zero. The universal finite identity reducing (II.180.2) to (II.180.9) is certified by `code/verify.py` and `data/grid_certificate.js on`.

Proof and precise finite component. Substitute (II.179.6) into (II.180.2). For each $T \subseteq \mathcal{G}$, its rational coefficient is

$$\alpha_T = \sum_{\substack{S \in \mathcal{S} \\ T \text{ fortyfour} \subseteq S}} (-1)^{|S| - |T|} \frac{\nu(S) - 2}{2\nu(S)c_1(S)c_2(S)}. \quad (\text{II.180.10})$$

If two finite step sets are related by an invertible rational affine map, their energies are equal: a pair of d -term sums is equal before the map if and only if it is equal afterward. The common affine translation cancels because both sums have length d .

The finite certificate evaluates (II.180.10) on all 512 supports, not on a finite list of values of d . To classify rational affine shapes, map every ordered noncollinear triple of a planar support to $(0, 0), (1, 0), (0, 1)$ and choose the lexicographically least resulting set of rational coordinate pairs. The corresponding one-dimensional normalization uses ordered distinct pairs; empty sets and singletons are treated separately. Two supports have the same key exactly when they are rationally affinely equivalent. Summing the coefficients (II.180.10) by these keys gives the following complete nonzero list:

Shape	P_7	P_6	P_{5a}	P_{5b}	P_4	P_3	P_2	P_1	$\emptyset$
Coefficient	1/12	1/6	-1/2	1/8	-1/2	7/6	-1	19/24	-1/3

This table is the finite computational component of the theorem. Its checker reconstructs all template directions from the 27 affine assignments, computes every rational coefficient, and verifies equality of the full dictionaries of affine-shape coefficients. All zero shape coefficients are checked as well. The empty-set energy is zero when $d \geq 1$, yielding (II.180.9). $\square$

[write] **7 October 2026; sources 40–46.** At $d = 8$: $H_8 = 2,585,444$ is source 18's C_8 , $2^{16} - 1 - H_8 = -2,519,909$ (checked at the write), and the second coefficient of the degenerate count is $\binom{k}{2} - 2,519,909k - b_{2,8}$, $b_{2,8} = 3,280,234,631,541$ (the certificate's output, reproduced on a copy at the write). At $d = 2$ the formula gives source 27's $(k^2 + 21k - 6)/2$.

The general fixed-grid theorem does not depend on this compression or its finite certificate. Thus the main algorithm and its proof survive even if the eight-shape identity is not used. The certificate is not a proof-assistant theorem: it is an explicit finite exact-arithmetic component of an otherwise conventional proof.

In the implementation the grid is ordered lexicographically and a support is encoded by its nine-bit mask. Representative masks for the displayed shapes are 254, 95, 31, 124, 15, 11, 10, 8, respectively. The certificate includes the actual point sets, so this convention can be checked without trusting the names of the shapes. In total 458 supports span the plane. Their multiplicities are

$$\#\{S : \nu(S) = 2, 3, 4, 6\} = 178, 128, 76, 76, \quad (\text{II.180.11})$$

and no other multiplicity occurs in the enumeration.

II.180.2 Exact values and independent cross-checks

The recurrences (II.179.5) and the general formula produce the following values. The first three nontrivial rows were also checked by independent enumeration of rational hyperplane normals and all pair spans, using primitive exterior-product coordinates rather than grid supports.

d	H_d	$b_{2,d}$	$\binom{H_d}{2} - b_{2,d}$
2	4	3	3
3	51	889	386
4	484	97 081	19 805
5	4 225	8 063 205	859 995
6	35 971	611 183 390	35 755 045
7	304 682	44 934 144 969	1 481 263 252
8	2 585 444	3 280 234 631 541	62 024 414 305

The values for $d = 9, 10$, together with every intermediate count and the multiplicity histograms, are included in `data/verification.json`. The table entries are outputs of a proved formula with exact integer recurrences; the independent small- d enumeration is an additional check, not the sole justification for extrapolating to $d = 8$.

For $d = 8$, the multiplicity counts are

$$\begin{array}{c|c}
 \nu & f_{\nu,8} \\
 \hline
 2 & 3\,157\,571\,257\,708 \\
 3 & 60\,640\,725\,488 \\
 4 & 460\,052\,299 \\
 6 & 353\,192
 \end{array} \tag{II.180.12}$$

They satisfy both identities (II.179.9). In particular,

$$\binom{2\,585\,444}{2} = 3\,342\,259\,045\,846, \quad b_{2,8} = 3\,280\,234\,631\,541.$$

Since $2^{16} - 1 - H_8 = -2\,519\,909$, Corollary II.178.3 gives (F.49.2). The threshold $p > 16$ is enough for every fixed k : it is not necessary to impose $p > 16 \cdot 2^k$.

The first newly evaluated instance is already informative. For $d = 3$,

$$\Delta_{k,3}(q) = \frac{k+51}{q} + \frac{\binom{k}{2} + 12k - 889}{q^2} + O_k(q^{-3}) \quad (p > 16). \tag{II.180.13}$$

For $d = 2$, the same formula reduces to the previously known coefficient $(k^2 + 21k - 6)/2$ [200]. This agreement tests the normalization of $b_{2,d}$ and the sign of its contribution.

II.181 The exponentially smaller pair-overlap deficit

Source 44, delivered Section 9.

The large count $b_{2,d}$ is close to the number of hyperplane pairs. The difference has its own growth scale. We derive that scale directly from the fixed-grid theorem, so the following asymptotic does not depend on the eight-shape finite certificate.

Lemma II.181.1 (The dominant seven-point support). *For every $S \in \mathcal{S}$, $\nu(S) \leq 6$. If $\nu(S) \geq 3$, then $|S| \leq 7$. Among seven-point supports there are exactly two with $\nu(S) \geq 3$: the central strips $|X + Y| \leq 1$ and $|X - Y| \leq 1$ inside $\mathcal{G}$. Each has $\nu(S) = 3$ and $c_1(S) = c_2(S) = 2$.*

Proof. Restrict any admissible nonconstant affine function to three noncollinear points of S . Modulo an affine change of its values, a nonconstant ternary profile on three labelled points has six possibilities: three bipartitions for two-valued profiles, and three orderings modulo reversal for three-valued profiles. Distinct gradient directions give distinct classes upon restriction, so $\nu(S) \leq 6$.

The two coordinate projections already give two directions. A third requires a function with both gradient coordinates nonzero. If $|S| \geq 8$, at least two horizontal rows are full. The three values on one full row are exactly $\{-1, 0, 1\}$, and the values on the other are a nonzero translate of that set. A finite nonempty subset of $\mathbb{Q}$ is not invariant under a nonzero translation. Thus $|S| \geq 8$ is impossible.

If $|S| = 7$, at least one row and one column are full. Their three-term arithmetic progressions force both gradient coordinates to be 1 or -1 , after the normalization that the output alphabet

is $\{-1, 0, 1\}$. The allowed grid points lie in a strip of three consecutive levels of $X + Y$ or $X - Y$. A noncentral such strip has at most six grid points, whereas the central strip has exactly seven. These are therefore the only possibilities. On either central strip, the same full row and column argument shows that the only extra non-axis direction is its defining diagonal; the other diagonal takes more than three values. This proves $\nu = 3$. Both coordinate projections have three values, so $c_1 = c_2 = 2$. $\square$

The two seven-point supports are equivalent by reflection to the support of P_7 . Their combined weight in (II.180.2) is $1/12$. Every remaining contributing support has at most six points. Since $0 \leq N_d(S) \leq |S|^{2d}$ and all weights in (II.180.2) are at most $1/12$, the crude but explicit bound

$$\left| B_d - \frac{1}{12} \mathcal{E}_d(P_7) \right| \leq 44 \cdot 36^d \quad (\text{II.181.1})$$

follows. Indeed there are at most 512 remaining supports, and the sequences on the seven-point support missing a specified point number at most 6^{2d} . Summing the seven missing-point bounds contributes at most $(7/12)36^d$; $(512 + 7)/12 < 44$.

Theorem II.181.2 (Triangular-lattice deficit asymptotic). *As $d \rightarrow \infty$,*

$$B_d = \frac{7 \cdot 49^d}{96\pi\sqrt{3}d} \left(1 - \frac{5}{24d} + O(d^{-2}) \right). \quad (\text{II.181.2})$$

Furthermore,

$$b_{2,d} \sim \frac{3 \cdot 81^d}{64\pi d}, \quad \frac{B_d}{\binom{H_d}{2}} \sim \frac{14}{9\sqrt{3}} \left(\frac{49}{81} \right)^d. \quad (\text{II.181.3})$$

Proof. The support of P_7 is centrally symmetric, so $\mathcal{E}_d(P_7) = \text{CT } P_7^{2d}$. Fourier coefficient extraction gives

$$\mathcal{E}_d(P_7) = \frac{1}{(2\pi)^2} \int_{[-\pi, \pi]^2} (1 + 2 \cos u + 2 \cos v + 2 \cos(u - v))^{2d} du dv. \quad (\text{II.181.4})$$

The absolute value of the bracket is at most seven, with equality only at $(0, 0)$ on the torus. This follows from the triangle inequality for the seven unit complex terms: equality requires all of them to have the same phase, which the constant term forces to be one.

Put $Q_2 = u^2 + v^2 + (u - v)^2$ and $Q_4 = u^4 + v^4 + (u - v)^4$. Near zero, the logarithm of the bracket divided by seven is

$$-\frac{Q_2}{7} + \frac{Q_4}{84} - \frac{Q_2^2}{98} + O((u^2 + v^2)^3). \quad (\text{II.181.5})$$

After $(u, v) = \sqrt{7/(2d)}(s, t)$, the exponent $2d$ times this logarithm is

$$-Q_2(s, t) + \frac{1}{d} \left(\frac{7}{24} Q_4(s, t) - \frac{1}{4} Q_2(s, t)^2 \right) + O(d^{-2}(s^2 + t^2)^3).$$

The Gaussian integral is $\int_{\mathbb{R}^2} e^{-Q_2} ds dt = \pi/\sqrt{3}$. Under the normalized Gaussian its covariance matrix is

$$\begin{pmatrix} 1/3 & 1/6 \\ 1/6 & 1/3 \end{pmatrix}.$$

Hence $\mathbb{E}Q_4 = 1$ and $\mathbb{E}Q_2^2 = 2$. For example each of $s, t, s - t$ has variance $1/3$ and fourth moment $1/3$; also Q_2 is a sum of two squares after linear normalization, giving the stated second moment. The relative coefficient at order $1/d$ is therefore $7/24 - 1/2 = -5/24$.

Here is a remainder justification for this use of the expansion. Outside a fixed small neighborhood of zero the integrand is bounded by $(7 - \epsilon)^{2d}$ for some $\epsilon > 0$. Inside that neighborhood the real analytic logarithm has the displayed Taylor expansion and its quadratic form is positive definite. Restricting the scaled variables first to a ball of radius $d^{1/20}$ makes the integrated remainder, including the square of the fourth-order correction, bounded by d^{-2} times an integrable Gaussian times a fixed polynomial. The remaining scaled annulus is exponentially small in $d^{1/10}$, by the quadratic upper bound on the logarithm. Gaussian tails permit extension to $\mathbb{R}^2$. These bounds justify an $O(d^{-2})$ relative error. Thus

$$\mathcal{E}_d(P_7) = \frac{7 \cdot 49^d}{8\pi\sqrt{3}d} \left(1 - \frac{5}{24d} + O(d^{-2})\right).$$

Combine this with (II.181.1); its exponential error is smaller than every inverse power correction on the scale $49^d/d$.

For the final assertions, the one-dimensional coefficient integral for T_{2d} has integrand $(1 + 2 \cos u)^{2d}$. Its unique absolute maximum is three at zero, with normalized logarithm $-u^2/3 + O(u^4)$. The same Gaussian argument gives $T_{2d} \sim 9^d \sqrt{3}/(2\sqrt{2\pi d})$. The term $\binom{2d}{d} \leq 4^d$ is exponentially smaller, so $H_d \sim 9^d \sqrt{3}/(4\sqrt{2\pi d})$. Squaring and dividing by two gives $\binom{H_d}{2} \sim 3 \cdot 81^d/(64\pi d)$. The deficit is exponentially smaller, giving both statements in (II.181.3). $\square$

The dominant step set is a triangular-lattice hexagon with a holding step at its center. Its appearance is forced by the extremal grid support argument, not introduced as an analogy after the asymptotic was observed numerically. Further coefficients can be obtained from the same analytic integral, but their all-orders derivation is not needed here.

II.182 Contrasting growing-dimension laws

Source 44, delivered Section 10.

The exact products allow regimes inaccessible to a fixed-order inclusion–exclusion approximation. Their interpretation must remain precise: Φ is the probability of avoiding the classified obstructions, not necessarily of avoiding every ternary relation.

Theorem II.182.1 (Signed-collision scale). *Fix $d \geq 2$. Along any sequence of finite fields with $p \geq 5$ and $q \rightarrow \infty$, if $k/q \rightarrow c \in [0, \infty)$, then*

$$\Phi_{k,d,p}(q) \longrightarrow e^{-c-c^2}. \quad (\text{II.182.1})$$

Consequently the probability of full nondegeneracy has upper limit at most e^{-c-c^2} in this regime.

Proof. The number of cross-section hyperplanes is bounded in terms of d alone, even when the characteristic varies, so $g_{d,p}(q) = 1 + O_d(q^{-1})$. The nonzero-side probability converges to e^{-c} . With $L = 2^m - 1$,

$$\log F_{k,d,p}(q) = -\frac{kL + k(k-1)}{q^2} + O_d\left(\frac{k + k^3}{q^4}\right) \longrightarrow -c^2.$$

All factors are positive for large q in this regime. Multiplying proves (II.182.1). Full nondegeneracy implies avoidance of the classified obstructions, so its probability is at most Φ . $\square$

The linear term c is the zero-side cost; the quadratic term c^2 is the signed-collision cost. The finite binary offset constellation has fixed size and does not contribute at this scale.

Theorem II.182.2 (Small-characteristic rank window). *Fix $p \in \{2, 3\}$, $d \geq 2$, and an integer $s \geq 0$. For $q = p^a$ and $k = 2a - m + 1 - s$, as $a \rightarrow \infty$,*

$$\Phi_{k,d,p}(q) \longrightarrow \prod_{j=s+1}^{\infty} (1 - p^{-j}). \quad (\text{II.182.2})$$

If $k > 2a - m + 1$, every arrangement is degenerate.

Proof. The last assertion is Corollary II.176.4. In the specified window the exact core product is

$$\prod_{i=0}^{k-1} (1 - p^{m-1+i-2a}) = \prod_{j=s+1}^{2a-m+1} (1 - p^{-j}).$$

Its limit is the displayed convergent positive product, because $\sum_j p^{-j} < \infty$. The cross-section product in (II.176.7) and the nonzero-side probability tend to one; the latter uses $k = O(a)$ and $a/p^a \rightarrow 0$. $\square$

Thus the classified obstruction has a dimension scale of order q in characteristics at least five, but a rank threshold of order $\log q$ in characteristics two and three. Theorem II.177.2 does not turn these into full nondegeneracy limits: its inventory $R_{k,d,p}$ grows rapidly with k . Any such strengthening would require new control of non-core directions, not simply another expansion of the exact product.

II.183 Source 45's chapter: energy levels and the coset range [write]

Added in the eighth write, 7 October 2026.

The next two sections are source 45's Sections 2–3. *Logically they follow source 14's near-maximal energy material* (Theorem II.43.2 and the collision results of its Section II.43) and source 15's collision amplification (Section II.45); they are printed here because no number printed earlier may move. Source 45's introduction is Section F.50; its Sections 4–5 are in Part III (Section III.89).

What is proved. For $m \geq 5$ the three largest additive energies of m -element subsets of a torsion-free abelian group are M_m , $M_m - 4(m - 2)$ and $M_m - 8(m - 3)$, $M_m = (2m^3 + m)/3$; the sets at the second value are the affine images of $H_m = \{0, \dots, m - 2, m\}$ (Theorems II.184.1, II.184.10, II.184.11), through the exact identity $(M_m - E(A))/4 = \#\{x < y < z : x + z - y \notin A\}$ (Proposition II.184.2) and a unit-defect theorem for positive Schur triples. The first gap follows from the inverse Pollard theorem and Hegyvári's work, as source 45 says (Remark II.184.7); the classification and the third gap are new here, and the third level is not classified. Source 14's envelope $|A\Delta C| \leq \tau(\varepsilon)m$, proved for $\varepsilon \leq 1/100$, is extended to $\varepsilon < 1/9$ in every abelian group, with the canonical subgroup $\{d : |A \cap (A + d)| > m/2\}$, uniqueness, and all positive equality cases, which are exactly the punctured cosets $C \setminus K$, K a subcoset of index at least nine (Theorems II.185.4, II.185.8); a second, nested-hole envelope with optimal leading constant one (Theorem II.185.11); for coordinate-cube collisions, equality forces one punctured coordinate (Theorems II.185.14, II.185.15). The equality cases are the intake's punctured-coset observation recorded at item II.31, now with all equality cases on the larger range. The largest range is open.

Relation to the formal statements (checked at HEAD 6d4942222). `theorem_7_1` and `theorem_7_2` (Sections06_07.lean, lines 104 and 125) have, since the seventh write's pins, the exact companions `theorem_7_1_holds` and `theorem_7_2_holds` (Proofs07FreimanClosure.lean, lines 185 and 235); the exact energy levels and the coset envelope are different statements and are not formalized.

II.184 The first three additive-energy levels

Source 45, delivered Section 2.

The opening discussion of Section 7 of Gowers's paper records an exact extremal fact: an arithmetic progression maximizes the additive energy of a fixed number of integers, and it is the only maximizer [1, Section 7, immediately before Theorem 7.2]. Here we examine the next two values below that maximum. This is a different endpoint from the subgroup-coset bound $E(A) \leq |A|^3$ used elsewhere in this article. The ordered setting has maximum $(2|A|^3 + |A|)/3$, and its first two positive defects have exact linear sizes.

The maximum, an endpoint-addition recurrence, and the congruence of energy modulo four belong to the existing literature. Hegyvári explicitly studies the set of attainable energies and computes the displaced-endpoint examples used below [189, Section 2.1]. The first gap also follows from the inverse Pollard theorem [190]; we give both its provenance and an elementary proof. The full second-level equality classification and the exact third level are extensions developed here beyond the statements located in the inspected sources and the pinned repository [202]. We make no assertion of global publication priority for them. Their proofs are self-contained.

II.184.1 Statement of the spectrum theorem

For a nonempty finite set $A \subset \mathbb{R}$, put

$$E(A) = \#\{(a, b, c, d) \in A^4 : a + b = c + d\}, \quad (\text{II.184.1})$$

$$M_m = \frac{2m^3 + m}{3}, \quad D(A) = \frac{M_m - E(A)}{4}, \quad m = |A|. \quad (\text{II.184.2})$$

The notation $D(A)$ in this section denotes a quarter of the deficit from the torsion-free maximum. Its integrality and nonnegativity will follow from a counting identity, rather than being assumed.

For $m \geq 3$, define

$$H_m = \{0, 1, \dots, m-2, m\}, \quad L_m = \{0, 2, 3, \dots, m\} = m - H_m. \quad (\text{II.184.3})$$

An affine image of a finite real set means its image under $x \mapsto u + vx$ with $v \neq 0$. In particular, reflection is allowed.

Theorem II.184.1 (The first three levels). *For every integer $m \geq 5$, the three largest distinct additive energies of m -element subsets of $\mathbb{R}$ are*

$$\boxed{M_m, \quad M_m - 4(m-2), \quad M_m - 8(m-3)}. \quad (\text{II.184.4})$$

The maximum occurs precisely for arithmetic progressions. The second value occurs precisely for affine images of H_m . If A belongs to neither of these two classes, then

$$E(A) \leq M_m - 8(m-3), \quad (\text{II.184.5})$$

and equality is attained by $\{0, 1, \dots, m\} \setminus \{2\}$. The same assertions hold for finite subsets of every torsion-free abelian group, with the affine models interpreted using a nonzero group element as their common step.

[write] **7 October 2026; sources 40–46.** Checked at the write by exhaustion over the m -subsets of $\{0, \dots, n\}$: the three largest energies are 85, 73, 69 for $m = 5$ ($n = 12$) and 146, 130, 122 for $m = 6$ ($n = 13$), as the formula gives; the intake's search through $m = 8$ agrees and finds the third value also at shapes other than J_m (for example $\{0, 2, 3, 4, 6\}$ at $m = 5$), consistent with the third level being unclassified.

The theorem classifies every set at the second level; it does not classify all sets at the third level. The restriction $m \geq 5$ is necessary for the stated second-level classification. Small-cardinality exceptions are discussed in Subsection II.184.6.

II.184.2 An exact ordered-completion deficit

The following identity turns the energy deficit into a count of specific failed completions. It is also a convenient exact certificate for finite computations.

Proposition II.184.2 (Ordered-completion identity). *For every nonempty finite $A \subset \mathbb{R}$,*

$$D(A) = \#\{(x, y, z) \in A^3 : x < y < z, x + z - y \notin A\}. \quad (\text{II.184.6})$$

In particular $D(A)$ is a nonnegative integer and $E(A) \equiv |A| \pmod{4}$.

Proof. Write $A = \{a_1 < \dots < a_m\}$. For $m \geq 2$, let $B = A \setminus \{a_m\}$, $a = a_m$, and $n = m - 1$. Define

$$T(B, a) = \#\{(x, y, z) \in B^3 : x + y - z = a\}, \quad (\text{II.184.7})$$

$$e(B, a) = \binom{n}{2} - T(B, a). \quad (\text{II.184.8})$$

In every triple counted by $T(B, a)$ one has $x > z$, because $x - z = a - y > 0$. Conversely, a pair $z < x$ in B determines the only possible value $y = a - x + z$. Consequently

$$e(B, a) = \#\{(z, x) \in B^2 : z < x, a - x + z \notin B\} \geq 0. \quad (\text{II.184.9})$$

Count the additive quadruples in A by their occurrences of a . Quadruples with no occurrence contribute $E(B)$. Exactly one occurrence contributes $4T(B, a)$, one copy for each position. If there are exactly two occurrences, one must be on each side of the equation: placing both on the same side would require two smaller numbers to sum to $2a$. The remaining values must agree, giving $4n$ quadruples. Exactly three occurrences are impossible, and four occurrences contribute one. Therefore

$$E(A) = E(B) + 4m - 3 + 4T(B, a). \quad (\text{II.184.10})$$

This endpoint recurrence is also present, in difference-multiplicity notation, in Hegyvári's proof of the congruence modulo four [189, proof of Theorem 2.1].

The elementary identity

$$M_m - M_{m-1} = 4m - 3 + 4\binom{m-1}{2}$$

now gives

$$D(A) = D(B) + e(B, a). \quad (\text{II.184.11})$$

Starting with a singleton, whose deficit is zero, and iterating proves (II.184.6). Indeed, for $z < x < a$, the completion $a - x + z$ lies strictly between z and a . Its membership in A is therefore equivalent to membership in the prefix preceding a . Finally $M_m \equiv m \pmod{4}$, since $M_m - m = 2m(m - 1)(m + 1)/3$ is divisible by four. The congruence follows. $\square$

The recurrence admits a useful equality statement.

Lemma II.184.3 (Zero endpoint defect). *Let $B \subset \mathbb{R}$ be finite and nonempty, and let $a > \max B$. Then $e(B, a) = 0$ if and only if $B \cup \{a\}$ is an arithmetic progression.*

Proof. The assertion is immediate for $|B| = 1$, so assume $|B| = n \geq 2$ and translate $\min B$ to zero. Let

$$D^+(B) = \{x - z : x, z \in B, x > z\}.$$

If $e(B, a) = 0$, then $a - d \in B \setminus \{0\}$ for every $d \in D^+(B)$. The exclusion of zero follows from $a - d \geq a - \max B > 0$. Hence $|D^+(B)| \leq n - 1$. The reverse inequality follows from $B \setminus \{0\} \subset D^+(B)$, and thus $D^+(B) = B \setminus \{0\}$.

Let d be the least positive element of B . Whenever $b \in B$ and $b > d$, the difference $b - d$ is again in B . Repeated subtraction must terminate, and a positive terminal value below d is impossible. Thus all elements are integer multiples of d , and all intermediate multiples up to the largest are present. Therefore $B = \{0, d, \dots, (n - 1)d\}$. The containment $a - D^+(B) \subset B \setminus \{0\}$ is between sets of the same size, so it is equality. Comparing their least elements gives $a = nd$. The converse follows directly by completing pairs in a progression. $\square$

Corollary II.184.4 (The classical maximum). *For every nonempty finite real set A , one has $E(A) \leq M_{|A|}$, with equality if and only if A is an arithmetic progression.*

Proof. The inequality follows from Proposition II.184.2. If $|A| \geq 2$ and equality holds, then both nonnegative terms in (II.184.11) vanish, and Lemma II.184.3 makes A a progression. Singletons are progressions as well. Conversely every completion in (II.184.6) succeeds for a progression. $\square$

II.184.3 Appending a point and the sharp first gap

Lemma II.184.5 (Appending a point to a progression). *Let $B = \{0, d, \dots, (n - 1)d\}$, where $d > 0$ and $n \geq 2$, and let $a > (n - 1)d$. If a/d is not an integer, then $T(B, a) = 0$. If $a = qd$ for an integer $q \geq n$, then*

$$T(B, a) = \sum_{h=\max\{1, q-n+1\}}^{n-1} (n - h), \quad (\text{II.184.12})$$

where an empty sum is zero. Consequently, if $B \cup \{a\}$ is not a progression, then $e(B, a) \geq n - 1$. For $n \geq 3$, equality in this last bound occurs precisely when $a = (n + 1)d$. If $n \geq 4$ and a is neither nd nor $(n + 1)d$, then

$$e(B, a) \geq 2n - 3. \quad (\text{II.184.13})$$

Proof. All values of $x + y - z$ on B^3 are integer multiples of d . For $a = qd$, a positive difference $hd = x - z$ has exactly $n - h$ representations in B , and its associated completion belongs to B exactly when $0 \leq q - h \leq n - 1$. This proves (II.184.12).

At $q = n$ the sum is $\binom{n}{2}$. For $q \geq n + 1$ it is at most $\binom{n-1}{2}$, attained only at $q = n + 1$ when $n \geq 3$. If a/d is not an integer, then $e(B, a) = \binom{n}{2} \geq n - 1$; this inequality is strict for $n \geq 3$. For $q \geq n + 2$, the sum is at most $\binom{n-2}{2}$, so the deficit is at least $2n - 3$. The nonintegral case also satisfies this bound because $\binom{n}{2} \geq 2n - 3$ for $n \geq 4$. $\square$

Theorem II.184.6 (Sharp first gap). *If $A \subset \mathbb{R}$ has $m \geq 3$ elements and is not an arithmetic progression, then*

$$D(A) \geq m - 2, \quad E(A) \leq M_m - 4(m - 2). \quad (\text{II.184.14})$$

Equality is attained by H_m .

Proof. For $m = 3$, remove the largest point. The remaining two-point set is an AP, and Lemma II.184.5 gives $D(A) = 1$ unless A is an AP. Now induct on m . Write $B = A \setminus \{\max A\}$. If B is an AP, Lemma II.184.5 gives $D(A) = e(B, \max A) \geq m - 2$. Otherwise the induction hypothesis gives $D(B) \geq m - 3$, and Lemma II.184.3 gives $e(B, \max A) \geq 1$. Their sum is at least $m - 2$. For H_m , append m to $\{0, \dots, m - 2\}$ and use the equality case in Lemma II.184.5. $\square$

Remark II.184.7 (Why the first gap is a consequence of earlier work). We record the precise relationship to inverse Pollard, so that the elementary proof above is not mistaken for an assertion of novelty of the bound. For an integer set A of size m , let

$$r_A(d) = |A \cap (A + d)|, \quad S_t(A) = \sum_d \min\{t, r_A(d)\}.$$

Embed A in a sufficiently large prime cyclic group, with no new two-sum equalities. Pollard's inequality gives $S_t(A) \geq t(2m - t)$. For $2 \leq t \leq m - 2$, the equality classification in [190, Theorem 3] says that a non-AP A gives strict inequality: none of the exceptional endpoint cases applies to $(A, -A)$. At $t = 1$, use the classical equality case in $|A - A| \geq 2m - 1$, or the recalled Vosper theorem in [190, Theorem 1]. A cyclic AP in these equality cases pulls back to an integer AP, using its consecutive three-term relations.

Symmetry $r_A(d) = r_A(-d)$ and $r_A(0) = m$ show that $S_t(A)$ has parity t , as does $t(2m - t)$. Thus each of these $m - 2$ strict inequalities improves by at least two. The layer-cake identity

$$E(A) = (2m - 1)m^2 - 2 \sum_{t=1}^{m-1} S_t(A)$$

then gives $M_m - E(A) \geq 4(m - 2)$. This is the first-gap bound. Hegyvári's formula for $\{1, \dots, m - 1, m + k\}$ already gives the extremal example at $k = 1$ [189, Section 2.1].

For a finite set in a torsion-free abelian group, identify its generated subgroup with $\mathbb{Z}^r$ and choose an integer linear functional avoiding the finitely many nonzero differences of its pair sums. Its restriction is a Freiman 2-isomorphism to an integer set. It preserves energy and the non-AP condition: an AP in the image pulls back using consecutive three-term relations. Thus the same prior implication covers real sets and general torsion-free sets.

II.184.4 A unit deficit in the positive Schur-triple problem

The next classification provides information that the numerical first gap alone does not supply. Positivity is essential here; we count ordered Schur triples with repeated entries allowed.

Lemma II.184.8 (Exactly one missing Schur triple). *Let $C = \{c_1 < \cdots < c_n\} \subset (0, \infty)$, and put*

$$S(C) = \#\{(x, y, z) \in C^3 : x + y = z\}, \quad h(C) = \binom{n}{2} - S(C).$$

Then $h(C) \geq 0$, with equality precisely when $C = d\{1, 2, \dots, n\}$ for some $d > 0$. If $n \geq 4$, then

$$h(C) = 1 \iff C = d\{1, 2, \dots, n-1, n+1\} \quad (d > 0). \quad (\text{II.184.15})$$

Proof. Let r_j be the number of ordered representations $x + y = c_j$ in C . Both summands precede c_j , and each of the $j-1$ possible first summands determines the second. Therefore

$$h(C) = \sum_{j=1}^n (j-1-r_j), \quad j-1-r_j \geq 0. \quad (\text{II.184.16})$$

Equality $r_j = j-1$ says exactly that the prefix $\{c_1, \dots, c_{j-1}\}$ is invariant under the reflection $x \mapsto c_j - x$. If this holds at every index, induction gives $c_j = jd$, with $d = c_1$. This proves the zero-deficit assertion.

Suppose now that $h(C) = 1$. There is a unique deficient index k , whose contribution in (II.184.16) is one. Before that index, $c_i = id$. Necessarily $k \geq 2$.

First suppose $k \geq 3$. We have $r_k = k-2 \geq 1$, so $c_k = qd$ for an integer $q \geq k$. The case $q = k$ has zero deficit and is excluded. For $k \leq q \leq 2k-2$, the number of ordered pairs in $\{1, \dots, k-1\}^2$ summing to q is $2k-1-q$; above this range it is zero. The equality $r_k = k-2$ therefore forces $q = k+1$. If $k < n$, zero deficit at the next index would require the prefix

$$\{d, 2d, \dots, (k-1)d, (k+1)d\}$$

to be invariant under reflection about $c_{k+1}/2$. Comparing its least and greatest elements gives $c_{k+1} = (k+2)d$. But that reflection sends the present element $2d$ to the absent element kd , a contradiction. Thus $k = n$, giving the displayed model.

It remains to rule out $k = 2$. Here $c_2 \neq 2d$. Since $n \geq 4$, zero deficit at index three gives $c_3 = d + c_2$. Zero deficit at index four requires the three-point prefix $\{d, c_2, d + c_2\}$ to be centrally symmetric. Its middle point must be its center, so

$$2c_2 = d + (d + c_2),$$

forcing $c_2 = 2d$, a contradiction. Conversely, the model in (II.184.15) has $r_j = j-1$ for $j < n$ and $r_n = n-2$, and hence has deficit one. $\square$

Corollary II.184.9 (Unit endpoint defect). *Let $B \subset \mathbb{R}$ have $n \geq 4$ elements and $a > \max B$. If $e(B, a) = 1$, then $B \cup \{a\}$ is an affine image of L_{n+1} . More precisely, after translating $\min B$ to zero, there is a $d > 0$ such that*

$$a = (n+1)d, \quad B = \{0, 2d, 3d, \dots, nd\}.$$

Proof. Set $C = a - B \subset (0, \infty)$. The bijection $(x, y, z) \mapsto (a-x, a-y, a-z)$ sends the equation $x + y - z = a$ to the Schur equation $(a-x) + (a-y) = a-z$. Thus $e(B, a) = h(C)$. Apply Lemma II.184.8. Once $\min B = 0$, the maximum of C is a , so its classified form gives exactly the stated values. $\square$

II.184.5 The second and third levels

Theorem II.184.10 (All second-level sets). *Let $A \subset \mathbb{R}$ have $m \geq 5$ elements. Then*

$$D(A) = m - 2 \iff A \text{ is an affine image of } H_m.$$

Proof. The models attain the value by Theorem II.184.6. Conversely, write $B = A \setminus \{a\}$, with $a = \max A$. If B is an AP, the equality statement in Lemma II.184.5 gives an affine image of H_m . If B is not an AP, then Theorem II.184.6 and Lemma II.184.3 give

$$D(B) \geq m - 3, \quad e(B, a) \geq 1.$$

Their sum is $m - 2$, so $e(B, a) = 1$. Since $|B| = m - 1 \geq 4$, Corollary II.184.9 applies and makes A an affine image of L_m , hence a reflected affine image of H_m . $\square$

Theorem II.184.11 (The sharp third gap). *If $m \geq 5$ and an m -element real set A is neither an AP nor an affine image of H_m , then $D(A) \geq 2m - 6$.*

Proof. Call the APs and the affine images of H_m the first two classes. For $m = 5$, Theorems II.184.6 and II.184.10 give $D(A) > 3$ outside those classes. Integrality then gives $D(A) \geq 4 = 2m - 6$.

Assume $m \geq 6$, put $n = m - 1$, and again let $B = A \setminus \{a\}$, $a = \max A$. There are three cases.

If B is an AP, normalize it as $\{0, 1, \dots, n - 1\}$. The values $a = n$ and $a = n + 1$ would put A in the first two classes. All other values satisfy $e(B, a) \geq 2n - 3$ by Lemma II.184.5, which is stronger than the desired $2n - 4 = 2m - 6$.

If B is a second-level set, normalize it by a translation and a positive dilation to either

$$B = H_n = \{0, 1, \dots, n - 2, n\} \quad \text{or} \quad B = L_n = \{0, 2, 3, \dots, n\}.$$

In both cases $D(B) = n - 2$, and the difference one has multiplicity $n - 2$. If a is nonintegral, every positive-difference completion $a - h$ is nonintegral, so $e(B, a) = \binom{n}{2} \geq n - 2$. If a is an integer at least $n + 2$, all the difference-one completions exceed $\max B$, again giving $e(B, a) \geq n - 2$. For the remaining value $a = n + 1$, the model L_n would extend to L_{n+1} and is excluded by the hypothesis on A . In the model H_n , difference two has multiplicity $n - 2$ and completion $n - 1$, which is absent. Thus $e(B, a) \geq n - 2$ in every admissible case, and $D(A) \geq 2n - 4$.

Finally, suppose B belongs to neither of the first two classes. Induction gives $D(B) \geq 2n - 6$. Endpoint defect zero would force A to be an AP by Lemma II.184.3, and endpoint defect one would force A into the second class by Corollary II.184.9. Hence $e(B, a) \geq 2$, and $D(A) \geq (2n - 6) + 2 = 2m - 6$. $\square$

Sharpness follows from the next exact formula, which also exhibits more of the energy spectrum near an interval.

Proposition II.184.12 (Deleting one point from an interval). *Let $m \geq 2$ and $1 \leq j \leq m - 1$. Then*

$$D(\{0, 1, \dots, m\} \setminus \{j\}) = j(m - j) - \min\{j, m - j\}. \quad (\text{II.184.17})$$

Proof. Put $A = \{0, 1, \dots, m\} \setminus \{j\}$. For $x < y < z$ in A , the completion $x + z - y$ lies in $[0, m]$, so it fails to belong to A precisely when it equals j . Solving for the middle element gives $y = x + z - j$. The strict inequalities $x < y < z$ are equivalent to $x < j < z$. There are $j(m - j)$ choices

$$0 \leq x \leq j - 1, \quad j + 1 \leq z \leq m.$$

Their determined middle value belongs to A except when $y = j$, or equivalently $x + z = 2j$. The excluded pairs are exactly $(j - t, j + t)$ with $1 \leq t \leq \min\{j, m - j\}$. Apply Proposition II.184.2 to finish the count. $\square$

At $j = 1$ the formula gives $D = m - 2$. At $j = 2$ and $m \geq 5$ it gives $D = 2m - 6$. In the latter model the unique gap of size two is interior, with a unit gap on each side, whereas in H_m and L_m that gap is first or last. Thus it belongs to neither of the first two classes. Since $0 < m - 2 < 2m - 6$ for $m \geq 5$, the bounds and examples prove (II.184.4) over $\mathbb{R}$.

II.184.6 Transfer, exact rigidity, and limitations

Proposition II.184.13 (Torsion-free transfer). *Theorems II.184.1, II.184.6, II.184.10, and II.184.11 hold for finite subsets of every torsion-free abelian group.*

Proof. The subgroup generated by a finite set A is a finitely generated torsion-free abelian group, hence isomorphic to $\mathbb{Z}^r$ for some finite r . Choose r real numbers linearly independent over $\mathbb{Q}$. They define an injective additive map $\mathbb{Z}^r \rightarrow \mathbb{R}$. This map preserves all additive quadruples, and hence $E(A)$ and $D(A)$.

It also preserves the stated equality classes in both directions. For an image AP, the common real step is the image of the difference between two consecutive points. Its relations pull back by injectivity. For an image of H_m or L_m with $m \geq 5$, there are adjacent model points separated by the unit step. Their preimages again supply an actual group element as the step, and injectivity pulls back every listed model point. The real theorems therefore give the asserted group theorems. Conversely, every nontrivial torsion-free abelian group contains an infinite cyclic subgroup, in which all the displayed models can be realized. $\square$

Corollary II.184.14 (An exact rigidity window). *Let A have $m \geq 5$ elements in a torsion-free abelian group. If*

$$E(A) > M_m - 8(m - 3),$$

then A is an AP or an affine image of H_m . In the latter case, deleting one point leaves an $(m - 1)$ -term AP. The strict energy threshold cannot be weakened to a non-strict one.

Proof. Apply Theorem II.184.11; remove the isolated end point in the H_m model. The single-hole model with $j = 2$ attains the threshold and belongs to neither permitted class. $\square$

For $m = 3$, every non-AP triple has energy $15 = M_3 - 4$; the two consecutive gaps can have any unequal ratio. For $m = 4$, the second-level classification also has extra families. For example, $\{0, 1, 4, 5\}$ has energy $36 = M_4 - 8$, but is not an affine image of H_4 . These examples explain why the equality classification and the three-level theorem start at $m = 5$.

The torsion-free assumption must be retained. A finite subgroup of size m has energy m^3 , exceeding M_m when $m > 1$. For a subset of a finite cyclic group, the conclusions are available only after an explicit *rectification*: a Freiman isomorphism of order two to a torsion-free set, preserving both directions of all two-sum relations. Merely choosing integer representatives does not establish such an isomorphism, because modular wraparound may create additional quadruples. This distinction is particularly relevant to Gowers's argument, which passes between integers and cyclic groups.

Finally, the window here has absolute width $O(m)$ around a quantity of order m^3 . Its relative width is therefore $O(m^{-2})$. It supplies exact rigidity at a fine discrete scale; it does not supply a constant-error near-extremizer theorem or a better global bound in Szemerédi’s theorem. Shao’s work addresses a different structural scale, obtaining large pieces in progressions of controlled rank from normalized energy inequalities [191]. The two types of result are complementary: the present arguments resolve exact finite energy levels and equality geometry close to Gowers’s classical Section 7 endpoint.

II.185 A larger coset-stability range and the geometry of equality

Source 45, delivered Section 3.

The additive-energy endpoint relevant to Gowers’s inverse and arrangement arguments is a subgroup coset. Quantitative information near that endpoint is useful only when it controls both the ambient approximation and the geometry of its errors. Source 14 of the existing repository report proves the sharp quadratic coset envelope below for defect at most $1/100$ (its delivered Theorem 4.2); its Question 4 asks for a larger validity range and for all equality cases. Source 15 subsequently develops collision-sensitive arrangement amplification and equality theory, while retaining that small-defect range. We use those results as the local comparison, and credit their envelope and collision framework [202]. Their place in the original argument is the energy and arrangement machinery of Gowers’s proof [1].

This section extends the same envelope to defect strictly below $1/9$, constructs its subgroup directly from autocorrelation, and classifies every case of equality in that range. An exact remainder measures two separate obstructions to equality: points outside the closest coset, and the additive-energy defect of the missing points inside it. This also gives a quantitative approximation by a coset with a subcoset removed. Finally, for axis-cube collisions, equality forces the removed coset to occur in a single coordinate. The range $1/9$ is a proved sufficient range; its optimality is not asserted.

II.185.1 Autocorrelation and an intrinsic subgroup

Throughout this section, G is an abelian group and $A \subset G$ is finite and nonempty. The ambient group need not be finite. Write

$$m = |A|, \quad r(d) = |A \cap (A + d)|, \quad E(A) = \sum_{d \in G} r(d)^2, \quad \varepsilon = 1 - \frac{E(A)}{m^3}. \quad (\text{II.185.1})$$

Thus $E(A)$ is the number of ordered additive quadruples in A . All sums involving r have finite support. The number ε in (II.185.1) is the *actual* energy defect. An externally supplied upper bound for it will be denoted by ε_0 .

For a finite subgroup coset C of G , put $d_C(A) = |A \Delta C|$, where Δ denotes symmetric difference. We use the normalized distance

$$\delta(A) = \frac{1}{m} \min_C d_C(A). \quad (\text{II.185.2})$$

The minimum exists: singleton cosets provide a nonempty collection of finite, nonnegative integer distances. Only finite cosets need to be considered; an infinite coset has infinite symmetric difference with A . Set $E(\emptyset) = 0$ when an auxiliary missing set is empty.

The elementary identities

$$\sum_d r(d) = m^2, \quad m^3 - E(A) = \sum_{a,b \in A} (m - r(a - b)) \quad (\text{II.185.3})$$

will be used repeatedly. The second follows by grouping the ordered pairs (a, b) according to their difference. We also have

$$r(u - v) \geq r(u) + r(v) - m. \quad (\text{II.185.4})$$

Indeed, the subsets $A \cap (A + u)$ and $A \cap (A + v)$ of A have intersection of size at least $r(u) + r(v) - m$; translation of their intersection injects into a translate of the overlap counted by $r(u - v)$.

Lemma II.185.1 (A pointwise autocorrelation gap). *For every $d \in G$,*

$$\frac{r(d)}{m} \left(1 - \frac{r(d)}{m} \right) \leq 2\varepsilon. \quad (\text{II.185.5})$$

Proof. Let $I = \{x \in A : x + d \in A\}$, so $|I| = r(d)$. For $x \in I$ and $b \in A$, inequality (II.185.4), applied to $x - b$ and $x + d - b$, gives

$$m - r(d) \leq (m - r(x - b)) + (m - r(x + d - b)).$$

Sum over $I \times A$. The first sum on the right is bounded by $m^3 - E(A)$ because $I \subset A$, and the second has the same bound because $I + d \subset A$. Consequently

$$m r(d)(m - r(d)) \leq 2(m^3 - E(A)).$$

Division by m^3 proves the claim. □

Lemma II.185.2 (Canonical subgroup and preliminary approximation). *Suppose $0 \leq \varepsilon < 1/9$, and define*

$$a = \frac{1 - \sqrt{1 - 8\varepsilon}}{2}, \quad H = \{d \in G : r(d) > m/2\}. \quad (\text{II.185.6})$$

Then H is a finite subgroup. There is a unique H -coset C containing more than $5m/6$ points of A . If

$$t = |A \setminus C|, \quad u = |C \setminus A|, \quad x = \frac{t}{m}, \quad \delta = \frac{t + u}{m}, \quad h = |H|, \quad (\text{II.185.7})$$

then

$$0 \leq x \leq \frac{a}{2} < \frac{1}{6}, \quad \delta \leq \frac{a - 2ax + 2x^2}{1 - a} \leq \frac{a}{1 - a} < \frac{1}{2}. \quad (\text{II.185.8})$$

Proof. The parameter a lies in $[0, 1/3)$ and satisfies $\varepsilon = a(1 - a)/2$. Lemma II.185.1 gives the dichotomy

$$\frac{r(d)}{m} \leq a \quad \text{or} \quad \frac{r(d)}{m} \geq 1 - a \quad (d \in G). \quad (\text{II.185.9})$$

Thus H consists precisely of the differences on the upper branch. It contains zero and is finite, since it is contained in $A - A$. If $d, e \in H$, then

$$\frac{r(d - e)}{m} \geq 1 - 2a > a.$$

The dichotomy forces $d - e \in H$. A nonempty finite set closed under differences is a subgroup.

Let n_j be the sizes of the intersections of A with the distinct H -cosets, including only the nonempty intersections, and write $S = \sum_j n_j^2$. Counting pairs whose difference lies in H gives $S = \sum_{d \in H} r(d)$. On H , we use $r(d)^2 \leq mr(d)$; off H , the lower branch gives $r(d)^2 \leq amr(d)$. Hence

$$(1 - \varepsilon)m^3 = E(A) \leq mS + am(m^2 - S),$$

and therefore

$$\frac{S}{m^2} \geq 1 - \frac{\varepsilon}{1 - a} = 1 - \frac{a}{2}. \quad (\text{II.185.10})$$

Because $\sum_j n_j/m = 1$,

$$\max_j \frac{n_j}{m} \geq \sum_j \left(\frac{n_j}{m} \right)^2 = \frac{S}{m^2}.$$

Choose the maximizing coset C . It contains at least $(1 - a/2)m > 5m/6$ points of A , so it is unique, and $x \leq a/2$.

The upper branch on H and the concentration in C also give

$$h(1 - a)m \leq S \leq (m - t)^2 + t^2.$$

For the second inequality, the sum of the squares of the remaining coset occupancies is at most the square of their sum t . Since $h/m = 1 + \delta - 2x$, rearrangement yields

$$\delta \leq \frac{a - 2ax + 2x^2}{1 - a}.$$

Here $0 \leq x \leq a/2 \leq a$, so $-2ax + 2x^2 \leq 0$. Finally $a < 1/3$ gives $a/(1 - a) < 1/2$. $\square$

II.185.2 An exact remainder around a coset

The next estimate is independent of the preceding construction. It separates the points outside a proposed coset from the points missing inside it, and retains the energy of the latter set.

Proposition II.185.3 (The outlier and missing-set remainder). *Let $C = c + H$ be any finite subgroup coset, and put*

$$A_0 = A \cap C, \quad T = A \setminus C, \quad K = C \setminus A, \quad b = |A_0|, \quad t = |T|, \quad u = |K|.$$

With $x = t/m$ and $\delta = (t + u)/m$, one has

$$E(A) \leq E(A_0) + 6bt^2 + 5t^3, \quad (\text{II.185.11})$$

$$E(A_0) = b^3 - bu(b - u) - (u^3 - E(K)), \quad (\text{II.185.12})$$

and consequently

$$\varepsilon \geq \delta(1 - \delta) + x(2 + \delta^2 - (8 + \delta)x + 2x^2) + \frac{u^3 - E(K)}{m^3}. \quad (\text{II.185.13})$$

The coefficient of x in (II.185.13) is not asserted nonnegative for an arbitrary proposed coset.

Proof. An additive quadruple in A cannot have exactly one entry outside C : three entries in a subgroup coset force the fourth into that same coset. Quadruples with no entries in T contribute $E(A_0)$. For exactly two entries in T , there are six choices of their positions. After their values and one of the two inside values have been fixed, the final value is determined. These cases therefore contribute at most $6bt^2$. Exactly three entries in T contribute at most $4t^3$, and four entries contribute $E(T) \leq t^3$. This proves (II.185.11).

For (II.185.12), translate C to H and write $h = |H| = b + u$. Expanding the four factors of $\mathbf{1}_{A_0} = \mathbf{1}_H - \mathbf{1}_K$ in the additive-quadruple count gives

$$E(A_0) = h^3 - 4h^2u + 6hu^2 - 4u^3 + E(K).$$

For example, a term with two prescribed entries in K and two in H has hu^2 solutions: one H -entry is free and the last is determined in H . The other mixed terms are counted in the same way, regardless of the positions of the entries. Substitution of $h = b + u$ gives (II.185.12).

Finally combine the two estimates, use $b/m = 1 - x$ and $u/m = \delta - x$, and expand:

$$\begin{aligned} \frac{m^3 - E(A)}{m^3} &\geq 1 - (1 - x)^3 + (1 - x)(\delta - x)(1 - \delta) - 6(1 - x)x^2 - 5x^3 + \frac{u^3 - E(K)}{m^3} \\ &= \delta(1 - \delta) + x(2 + \delta^2 - (8 + \delta)x + 2x^2) + \frac{u^3 - E(K)}{m^3}. \end{aligned}$$

This is (II.185.13). □

II.185.3 The sharp envelope up to defect 1/9

For $0 \leq z \leq 1/4$, define

$$\tau(z) = \frac{1 - \sqrt{1 - 4z}}{2}, \quad \tau(z)(1 - \tau(z)) = z. \quad (\text{II.185.14})$$

Theorem II.185.4 (Expanded range of the quadratic coset envelope). *If $E(A) > 8m^3/9$, there is a unique finite subgroup coset C minimizing $|A \triangle C|$, and*

$$|A \triangle C| \leq \tau(\varepsilon)m. \quad (\text{II.185.15})$$

Its subgroup is the canonical set $H = \{d : r(d) > m/2\}$, and C is the unique H -coset containing more than half of A .

Proof. Use the coset and parameters of Lemma II.185.2. Define

$$F(\delta, x) = 2 + \delta^2 - (8 + \delta)x + 2x^2.$$

Since $x \leq \delta$ and $x < 1/6$, its derivative in δ is $2\delta - x \geq 0$, and therefore

$$F(\delta, x) \geq F(x, x) = 2 - 8x + 2x^2 > \frac{13}{18} > 0.$$

The last strict inequality follows because the quadratic decreases on $[0, 1/6]$ and takes the value $13/18$ at $1/6$. The missing-set defect $u^3 - E(K)$ is nonnegative. Hence Proposition II.185.3 gives $\delta(1 - \delta) \leq \varepsilon$. As $\delta < 1/2$, the increasing branch of this quadratic can be inverted, yielding $\delta \leq \tau(\varepsilon)$.

It remains to prove uniqueness among all finite subgroup cosets, not merely among H -cosets. Put $\tau = \tau(\varepsilon)$. We have

$$\tau < \tau(1/9) = \frac{3 - \sqrt{5}}{6} < \frac{1}{5}. \quad (\text{II.185.16})$$

Suppose that two cosets C, C' both have symmetric-difference distance at most τm from A . Then

$$|C \cap C'| \geq (1 - 2\tau)m, \quad |C|, |C'| \leq (1 + \tau)m.$$

Their intersection is larger than half of either coset, because $1 - 2\tau > (1 + \tau)/2$. A nonempty intersection of subgroup cosets is a coset of the intersection of their subgroups. Its index in either subgroup is a positive integer, so an intersection larger than half forces both indices to equal one. The subgroups, and then the cosets, coincide. Any coset closer to A than the constructed C would also satisfy the same τ -bound. This proves unique global minimality. The intrinsic description follows from Lemma II.185.2. $\square$

[write] 7 October 2026; sources 40–46. *Against source 14* (checked at the write). Source 14's Theorem II.43.2 proves the same envelope and uniqueness for $\varepsilon \leq 1/100$; here the range is $\varepsilon < 1/9$ ($\tau(1/9) = 0.1273$ against $\tau(1/100) = 0.0101$). Equality example recomputed exactly: $A = \mathbb{Z}/36\mathbb{Z} \setminus \{0, 9, 18, 27\}$ has $m = 32$, $E(A) = 29,184$, $\varepsilon = 7/64$ and $|A\Delta(\mathbb{Z}/36\mathbb{Z})|/m = 1/8 = \tau(7/64)$. The equality cases (Theorem II.185.8) are the punctured cosets of the intake observation at item II.31, which advances source 14's Question 4.

[write] 7 October 2026; sources 60–69. *Range extended by source 68* (Section II.296): the subgroup H and the envelope $\tau(\varepsilon)$ hold for every $\varepsilon \leq 1/4$, optimal even in cyclic groups, and the closest coset is unique for $\varepsilon < 2/9$, also optimal (its Theorem II.298.1, Corollary II.298.4, Proposition II.298.6); the equality cases are the punctured cosets of index at least four (its Theorem II.298.5).

Corollary II.185.5 (A supplied upper bound for the defect). *If $0 \leq \varepsilon_0 < 1/9$ and $E(A) \geq (1 - \varepsilon_0)m^3$, the same canonical closest coset satisfies*

$$|A\Delta C| \leq \tau(\varepsilon_0)m.$$

Proof. The actual defect satisfies $0 \leq \varepsilon \leq \varepsilon_0$. Apply Theorem II.185.4 and the monotonicity of τ . $\square$

Remark II.185.6 (What changes from the existing result). The function τ and its sharp small-defect interpretation are credited to source 14 [202]. The subgroup construction in Lemma II.185.2 replaces its preliminary small-difference-set argument, and makes the same envelope valid throughout $0 \leq \varepsilon < 1/9$. Neither this theorem nor the comparison claims that $1/9$ is the optimal endpoint. The statements also hold for finite sets in infinite abelian groups, since their canonical subgroup is contained in the finite set $A - A$.

II.185.4 Every exact envelope extremizer

From now on, C is the unique closest coset from Theorem II.185.4, and $\delta = \delta(A) = |A\Delta C|/m$. Keep x, t, u, K as in (II.185.7), and define the excess over the quadratic envelope by

$$\eta = \varepsilon - \delta(1 - \delta). \quad (\text{II.185.17})$$

Proposition II.185.7 (A structural remainder above the envelope). *If $\varepsilon < 1/9$, then*

$$\eta \geq (2 - 8\delta + 2\delta^2)x + \frac{u^3 - E(K)}{m^3} \geq x + \frac{u^3 - E(K)}{m^3}. \quad (\text{II.185.18})$$

Proof. For $0 \leq x \leq \delta < 1/5$, the polynomial $F(\delta, x)$ in the proof of Theorem II.185.4 decreases in x : its derivative is $-(8 + \delta) + 4x < 0$. Thus

$$F(\delta, x) \geq F(\delta, \delta) = 2 - 8\delta + 2\delta^2.$$

Insert this in (II.185.13). The last coefficient is itself decreasing on the relevant interval, and

$$2 - 8\tau(1/9) + 2\tau(1/9)^2 = \sqrt{5} - \frac{11}{9} > 1.$$

Since $\delta \leq \tau(\varepsilon) < \tau(1/9)$, the second inequality follows. $\square$

Theorem II.185.8 (All equality cases in the quadratic envelope). *Suppose $\varepsilon < 1/9$. Equality*

$$\delta(A) = \tau(\varepsilon) \quad (\text{II.185.19})$$

holds precisely in the following cases:

1. *A is a finite subgroup coset, and $\varepsilon = 0$;*
2. *$A = C \setminus K$, where $C = c + H$ is a finite subgroup coset, K is a coset of a subgroup $L \leq H$ contained in C , and $r = [H : L] \geq 9$.*

In the second case,

$$\delta(A) = \frac{1}{r-1}, \quad \varepsilon = \frac{r-2}{(r-1)^2}, \quad (\text{II.185.20})$$

and C is the unique closest coset.

Proof. Condition (II.185.19) is equivalent to $\eta = 0$. Proposition II.185.7 forces $x = 0$ and $E(K) = u^3$. Thus $A \subset C$. If $u = 0$, then $A = C$. If $u > 0$, maximal energy says that every ordered triple $x, y, z \in K$ satisfies $x - y + z \in K$. Translate one point of K to zero. The resulting finite set contains zero and is closed under negation and addition, hence is a subgroup L . Therefore K is a coset of L inside C . Writing $r = [H : L]$, we obtain $u = |L|$, $m = (r-1)|L|$, and $\delta = 1/(r-1)$. The strict inequality $\delta < \tau(1/9)$ implies $r \geq 9$.

Conversely, let $A = C \setminus K$ with the stated index. Since $E(K) = u^3$, identity (II.185.12) gives

$$m^3 - E(A) = mu(m - u).$$

Consequently $\varepsilon = \delta_0(1 - \delta_0)$ with $\delta_0 = u/m = 1/(r-1)$. For $r \geq 9$ this defect is less than $1/9$ and $\delta_0 < 1/2$, so $\delta_0 = \tau(\varepsilon)$. The displayed coset C satisfies the envelope bound with equality. The uniqueness part of Theorem II.185.4 identifies it with the closest coset, proving the converse and (II.185.20). $\square$

Corollary II.185.9 (The discrete equality parameters). *Across all finite abelian groups, the nonzero equality defects in the range $\varepsilon < 1/9$ are exactly*

$$\left\{ \frac{r-2}{(r-1)^2} : r \in \mathbb{N}, r \geq 9 \right\}.$$

In a fixed ambient group only indices of its finite subgroup inclusions can occur. In $\mathbb{F}_p^n$ the admissible indices are $r = p^j \geq 9$, with $1 \leq j \leq n$. In a prime cyclic group the nontrivial examples in this range are complements of one point, for $p \geq 11$.

Proof. Necessity follows from Theorem II.185.8. Every integer $r \geq 9$ is realized by a cyclic group H of order r with $L = \{0\}$. The remaining assertions follow from the possible subgroup indices in the indicated groups. $\square$

Remark II.185.10. The missing-subcoset examples already appear in a write note adjoining source 14's Question 4 [202]. Theorem II.185.8 proves that they exhaust equality in the enlarged range. In particular, sharpness of the function τ does not mean equality is attainable at every real defect parameter. Equality in Corollary II.185.5, with a supplied bound ε_0 , additionally forces the actual defect to equal ε_0 , by the strict monotonicity of τ .

II.185.5 Stability of the punctured-coset extremizers

The exact remainder controls the outliers and the missing-set defect through a common budget. Applying the coset envelope to the missing set, while keeping that budget intact, gives a second sharp envelope for repair. The relevant parameter is η/δ^3 , because the missing set has size approximately δm and energy scale $(\delta m)^3$.

Theorem II.185.11 (The nested-hole envelope and all equality cases). *Suppose $\varepsilon < 1/9$, $\delta = \delta(A) > 0$, and*

$$0 \leq s = \frac{\eta}{\delta^3} < \frac{1}{9}, \quad \eta = \varepsilon - \delta(1 - \delta). \quad (\text{II.185.21})$$

Let C be the canonical closest coset and $K = C \setminus A$. There is a unique closest subgroup coset D to K ; it is a nonempty proper subcoset of C . This same D minimizes the repair inside the fixed outer coset C . Write

$$\Delta_C(A) = \frac{1}{m} \min_{\substack{D' \subsetneq C \\ D' \text{ a nonempty subgroup coset}}} |A \Delta(C \setminus D')|.$$

If $s > 0$, then

$$\Delta_C(A) \leq \delta \tau(s) = \frac{\tau(s)}{s} \frac{\eta}{\delta^2} \leq (1 + 2s) \frac{\eta}{\delta^2}. \quad (\text{II.185.22})$$

Equality in the first bound holds precisely when

$$A = C \setminus (D \setminus J) = (C \setminus D) \cup J, \quad (\text{II.185.23})$$

where $J \subsetneq D \subsetneq C$ are nonempty subgroup cosets and the index of the subgroup underlying J in that underlying D is at least nine. The assertion fixes the canonical outer coset C . If $s = 0$, then $A = C \setminus D$ and $\Delta_C(A) = 0$.

Proof. Keep $x = |A \setminus C|/m$ and put $k = |K|/m = \delta - x$. The coupled remainder (II.185.18) gives $x \leq \eta = s\delta^3 < \delta$, so $k > 0$ and K is nonempty. Define its actual defect by $\zeta = 1 - E(K)/|K|^3$. The same remainder states that

$$\eta \geq x + k^3 \zeta. \quad (\text{II.185.24})$$

To control the full interval of possible outlier masses, define

$$z(y) = \frac{\eta - y}{(\delta - y)^3} \quad (0 \leq y \leq \eta).$$

Its denominator is positive, and

$$z'(y) = \frac{-\delta + 3\eta - 2y}{(\delta - y)^4} < 0,$$

where $3\eta = 3s\delta^3 < \delta/3$ uses $s < 1/9$ and $\delta < 1$. Thus $0 \leq z(y) \leq z(0) = s < 1/9$ throughout the interval. In particular (II.185.24) gives $0 \leq \zeta \leq z(x) < 1/9$.

Apply Theorem II.185.4 to K inside C , after translating C to its finite subgroup. There is a closest subcoset $D \subset C$ with

$$|K \triangle D| \leq \tau(\zeta)|K|. \quad (\text{II.185.25})$$

The ambient version of the same theorem makes D the unique closest coset even among those not contained in C . Moreover $|K| \leq \delta m < m/5$, so

$$|D| \leq (1 + \tau(\zeta))|K| < \frac{5}{4}|K| < \frac{m}{4}, \quad |C| \geq (1 - \delta)m > \frac{4m}{5}.$$

Hence D is proper. The outliers are disjoint from both sets inside C , giving

$$\Delta_C(A) = x + \frac{|K \triangle D|}{m} \leq x + k\tau(\zeta) \leq x + (\delta - x)\tau(z(x)). \quad (\text{II.185.26})$$

The equality on the left also proves the asserted minimization property of D .

Suppose $s > 0$ and define

$$F(y) = y + (\delta - y)\tau(z(y)) \quad (0 \leq y \leq \eta).$$

Write $k_y = \delta - y$ and $b_y = \tau(z(y))$. Then $0 \leq b_y < 1/2$ and $b_y(1 - b_y) = (\eta - y)/k_y^3$. Differentiating gives

$$\begin{aligned} F'(y) &= 1 - b_y + \frac{-k_y^{-2} + 3b_y(1 - b_y)}{1 - 2b_y} \\ &= \frac{1 - b_y^2 - k_y^{-2}}{1 - 2b_y} < 0, \end{aligned}$$

because $0 < k_y \leq \delta < \tau(1/9) < 1$. Thus (II.185.26) implies $\Delta_C(A) \leq F(x) \leq F(0) = \delta\tau(s)$. For the simpler estimate, put $b = \tau(s) < \tau(1/9) < 1/4$. Then

$$\frac{\tau(s)}{s} = \frac{1}{1 - b} = 1 + \frac{s}{(1 - b)^2} \leq 1 + \frac{16}{9}s \leq 1 + 2s.$$

When $s = 0$, Theorem II.185.8 gives the exact punctured-coset representation, with no division by s required.

Suppose next that $s > 0$ and equality holds in the first bound. Strict decrease of F forces $x = 0$, so $A \subset C$ and $k = \delta$. The exact complement identity now gives $\eta = \delta^3\zeta$, hence $\zeta = s$. Equality in (II.185.26) therefore requires

$$\frac{|K \triangle D|}{|K|} = \tau(\zeta).$$

The positive-defect equality classification in Theorem II.185.8, applied to K , says exactly that $K = D \setminus J$, where J is a nonempty subcoset of D and its underlying subgroup has index at least nine. This is (II.185.23). Conversely, under the theorem's hypotheses, such a representation makes

D the unique closest coset to K by the same inner equality classification. The complement identity gives $\zeta = s$, and

$$\Delta_C(A) = \frac{|K \triangle D|}{m} = \delta \tau(s),$$

proving sufficiency as well. No minimization over different outer cosets is asserted in this equality statement. $\square$

[write] **7 October 2026; sources 60–69.** *Range extended by source 68:* outer defect below $1/4$ and inner defect $s \leq 1/4$, with all equality cases and an oriented form (its Theorems II.299.4, II.299.5, II.299.8, Proposition II.299.11).

Proposition II.185.12 (Attainment and the sharp leading constant). *The first bound in (II.185.22) is attained by nested-hole examples for every pair of integers $r \geq 10$ and $u \geq 8$, at $s = (u - 1)/u^2$. These examples have minimum repair cost exactly one point, even if the outer coset is allowed to vary. Consequently the scale η/δ^2 cannot be replaced by a little-oh bound, and the optimal uniform leading constant in $(C_0 + o(1))\eta/\delta^2$, as $\eta/\delta^3 \rightarrow 0$, is $C_0 = 1$. The assertion of sharpness concerns these attained families; it does not assert attainment for every real parameter.*

Proof. Let the ambient group H be cyclic of order $r(u + 1)$, let $L \leq H$ have order $u + 1$, and put

$$K = L \setminus \{0\}, \quad A = H \setminus K, \quad m = (r - 1)(u + 1) + 1.$$

The difference multiplicities of K inside L are u at zero and $u - 1$ at every nonzero element. Therefore

$$E(K) = u^2 + u(u - 1)^2 = u^3 - u^2 + u.$$

The exact complement identity gives

$$\delta_0 = \frac{u}{m}, \quad \eta_0 = \frac{u^2 - u}{m^3}, \quad \varepsilon = \delta_0(1 - \delta_0) + \eta_0. \quad (\text{II.185.27})$$

Since

$$\delta_0 < \frac{1}{r - 1} \leq \frac{1}{9}, \quad \eta_0 < \delta_0^2,$$

we have $\varepsilon < \delta_0 < 1/9$. Also $\delta_0 < 1/2$ and $\varepsilon \geq \delta_0(1 - \delta_0)$, so the coset H satisfies $|A \triangle H|/m = \delta_0 \leq \tau(\varepsilon)$. Uniqueness in Theorem II.185.4 makes H the canonical closest coset. Hence $\delta = \delta_0$, $\eta = \eta_0$, and

$$s = \frac{\eta}{\delta^3} = \frac{u - 1}{u^2} \leq \frac{7}{64} < \frac{1}{9}, \quad \tau(s) = \frac{1}{u}.$$

The bound on s uses $u \geq 8$ and decrease of $(u - 1)/u^2$ for $u > 2$.

The punctured coset $H \setminus L$ differs from A by one point. There is no zero-edit representation $A = C' \setminus D'$ by nested subgroup cosets in this ambient group. Indeed $|A| > |H|/2$, so a coset C' containing A must equal H : every proper subgroup coset of the finite group H has size at most $|H|/2$. An exact representation would then require $D' = K$. But a subgroup coset contained in L has size dividing $|L| = u + 1$, whereas $|K| = u$ does not divide $u + 1$ for $u > 1$. Thus the minimum edit count is exactly one, and

$$\Delta_H(A) = \frac{1}{m} = \delta \tau(s).$$

This proves finite attainment with all the stated numerical conditions. Finally

$$\frac{\eta}{\delta^2} = \frac{1 - 1/u}{m}, \quad \frac{1/m}{\eta/\delta^2} = \frac{u}{u - 1} \longrightarrow 1, \quad s \longrightarrow 0 \quad (u \rightarrow \infty). \quad (\text{II.185.28})$$

The upper bound in Theorem II.185.11 has multiplier $\tau(s)/s = 1 + s + O(s^2)$, which tends to one. The attained family proves that no smaller uniform asymptotic leading constant is possible. $\square$

II.185.6 Axis-cube collisions and coordinate splitting

We now distinguish ordinary additive energy from the more restrictive axis-cube collision count used in the arrangement framework. Let

$$G = \prod_{i=1}^k G_i, \quad k \geq 1,$$

where the G_i are abelian groups. For $g \in G$, let $e_i(g)$ denote the element with i th coordinate g_i and all other coordinates zero. For $x, h \in G$, define the axis cube

$$\text{Cube}(x, h) = \{x + \omega \cdot h : \omega \in \{0, 1\}^k\}, \quad \omega \cdot h = (\omega_1 h_1, \dots, \omega_k h_k). \quad (\text{II.185.29})$$

Degenerate cubes are included. For a finite set $A \subset G$, let

$$\mathcal{Q}(A) = \#\{(x, z, h) : \text{Cube}(x, h), \text{Cube}(z, h) \subset A\}. \quad (\text{II.185.30})$$

This count is finite even if the ambient groups are infinite: the endpoints x and $x + h$ determine h from two elements of A .

For completeness, if $\phi : A \rightarrow W$ takes values in an abelian group, define on a cube contained in A

$$D_h \phi(x) = \sum_{\omega \in \{0, 1\}^k} (-1)^{|\omega|} \phi(x + \omega \cdot h).$$

Let $\mathcal{Q}_\phi(A)$ count the triples in (II.185.30) for which $D_h \phi(x) = D_h \phi(z)$. This is the label-respected collision count. Changing all signs by a common factor changes neither the equality condition nor the count.

Lemma II.185.13 (Collision and corner bounds). *For every finite nonempty A of size m ,*

$$\mathcal{Q}_\phi(A) \leq \mathcal{Q}(A) \leq E(A) \leq m^3. \quad (\text{II.185.31})$$

If

$$C_i(A) = \#\{(x, w) \in A^2 : x + e_i(w - x) \in A\},$$

then

$$\mathcal{Q}(A) \leq m C_i(A) \quad (1 \leq i \leq k). \quad (\text{II.185.32})$$

Proof. Send (x, z, h) to $(x, z, w = x + h)$. This map is injective. A valid cube pair has $x, z, w, z + w - x \in A$, so its image is an additive-energy triple. This proves the middle inequality in (II.185.31); the other inequalities are immediate. For the same image, the corner $x + e_i(w - x)$ belongs to A . For each admissible pair (x, w) there are at most m choices of z , which proves (II.185.32). $\square$

Theorem II.185.14 (Product-coset stability below defect $1/9$). *Suppose*

$$\mathcal{Q}(A) \geq (1 - \varepsilon_0) m^3, \quad 0 \leq \varepsilon_0 < 1/9.$$

Then the unique closest finite subgroup coset to A is a coordinate product

$$C = \prod_{i=1}^k (c_i + H_i), \quad H_i \leq G_i, \quad (\text{II.185.33})$$

and $|A \triangle C| \leq \tau(\varepsilon_0) m$. The conclusion also follows from the same lower bound on $\mathcal{Q}_\phi(A)$.

Proof. By (II.185.31) and Corollary II.185.5, the unique closest coset $C = c + H$ exists. Write $h = |H|$, $t = |A \setminus C|$, $x = t/m$, and $\delta = |A \Delta C|/m$. Then

$$x \leq \delta \leq \tau(\varepsilon_0) < 1/5, \quad h/m = 1 + \delta - 2x.$$

We prove $e_i H \subset H$ for each i .

If this inclusion failed, the homomorphism

$$F_i : H \times H \longrightarrow G, \quad F_i(y, w) = y + e_i(w - y),$$

would have image $H + e_i H$. Its image contains H by taking $y = w$, and contains $e_i H$ by taking $y = 0$; the reverse inclusion follows from the formula. Since $e_i H \not\subset H$, this image has at least $2h$ elements. Every nonempty fiber consequently has size at most $h/2$.

After translation by c , at most $mh/2$ pairs in C^2 have their mixed corner in A . At most $m^2 - (m - t)^2$ pairs in A^2 are outside C^2 . We infer

$$\begin{aligned} \frac{C_i(A)}{m^2} &\leq 1 - (1 - x)^2 + \frac{1 + \delta - 2x}{2} \\ &= \frac{1}{2} + \frac{\delta}{2} + x - x^2 \leq \frac{1}{2} + \frac{3\delta}{2} < \frac{4}{5}. \end{aligned}$$

But (II.185.32) and the hypothesis give $C_i(A)/m^2 \geq \mathcal{Q}(A)/m^3 > 8/9$, a contradiction. Thus every coordinate component of every element of H lies in H . If H_i is its projection to G_i , summing these coordinate components shows $H = \prod_i H_i$. This proves (II.185.33); the distance bound and uniqueness were already obtained. The label-respected case follows from $\mathcal{Q}_\phi \leq \mathcal{Q}$. $\square$

II.185.7 Positive-defect equality requires one punctured coordinate

The energy theorem allows a missing subgroup coset with an arbitrary orientation inside the closest coset. Collision equality is more rigid.

Theorem II.185.15 (All positive-defect collision-envelope extremizers). *Let*

$$\varepsilon_{\text{cube}} = 1 - \frac{\mathcal{Q}(A)}{m^3}, \quad 0 < \varepsilon_{\text{cube}} < 1/9.$$

Then

$$\delta(A) = \tau(\varepsilon_{\text{cube}}) \tag{II.185.34}$$

if and only if, after translation, there are finite subgroups $H_i \leq G_i$ and one coordinate j such that

$$A = (H_j \setminus K_j) \times \prod_{i \neq j} H_i, \tag{II.185.35}$$

where K_j is a coset of a subgroup $L_j \leq H_j$ and $[H_j : L_j] \geq 9$. Its unique closest coset is $C = \prod_i H_i$, translated back to the original coordinates. At zero defect, $\mathcal{Q}(A) = m^3$ holds precisely for coordinate product cosets.

Proof. Let $\varepsilon = 1 - E(A)/m^3$ be the actual additive-energy defect. If (II.185.34) holds, then

$$\delta(A) \leq \tau(\varepsilon) \leq \tau(\varepsilon_{\text{cube}}) = \delta(A).$$

Strict monotonicity of τ forces $\varepsilon = \varepsilon_{\text{cube}}$, so $E(A) = \mathcal{Q}(A)$. Theorem II.185.8 gives $A = C \setminus K$, where K is a subgroup coset inside C , and Theorem II.185.14 makes C a coordinate product.

We first show that A is closed under coordinate mixing. Suppose $x, w \in A$ and $v = x + e_i(w - x) \notin A$. Since $A \subset C$ and $|C| = m(1 + \delta) < 2m$, there is $z \in A$ with $z + w - x \in A$: the two relevant translates of A inside C have intersection of size at least $2m - |C| > 0$. Thus (x, z, w) is an additive-energy triple, but its uniquely corresponding cube pair, with side $w - x$, fails at the corner v . The injection in Lemma II.185.13 is therefore not surjective, contrary to $E(A) = \mathcal{Q}(A)$. Every such mixed corner must lie in A .

A nonempty set closed under replacing one coordinate by the corresponding coordinate of any other member is the Cartesian product of its coordinate projections. To verify this, start with any member and successively replace coordinates to obtain any prescribed tuple from those projections. After translating C to $H = \prod_i H_i$, write

$$A = \prod_i A_i, \quad A_i \subset H_i, \quad K = b + L = H \setminus A.$$

All A_i are nonempty, and at least one is proper because the defect is positive. Suppose two coordinates have proper A_i . For every coordinate i , choose a different coordinate j with $A_j \neq H_j$, and fix its value in $H_j \setminus A_j$. The entire i th-coordinate fiber, with the other coordinates fixed arbitrarily, then lies in K . Differences within this fiber imply $e_i H_i \subset L$. This holds for every i , so $L = H$ and $K = H$, contradicting $A \neq \emptyset$. Exactly one coordinate, say j , is proper. Therefore

$$K = (H_j \setminus A_j) \times \prod_{i \neq j} H_i.$$

Projection of the coset K shows that $H_j \setminus A_j$ is a subgroup coset K_j . This gives (II.185.35), and its index is the index $[H : L] \geq 9$ from Theorem II.185.8.

Conversely, let A have the form (II.185.35). Both additive energy and axis-cube collisions factor over Cartesian products. In a full subgroup coordinate, every triple of the two bases and the side is valid, giving $|H_i|^3$. In the sole punctured coordinate, a one-dimensional cube pair is exactly an additive quadruple, so its collision count equals its energy. Hence $\mathcal{Q}(A) = E(A)$. Theorem II.185.8 proves (II.185.34) and identifies the closest coset.

At zero defect, (II.185.31) and Theorem II.185.14 force A itself to be a coordinate product coset. Every such coset has $\mathcal{Q}(A) = m^3$, proving the last assertion. $\square$

Corollary II.185.16 (What equality says for labels). *If $\varepsilon_\phi = 1 - \mathcal{Q}_\phi(A)/m^3 < 1/9$ and $\delta(A) = \tau(\varepsilon_\phi)$, then*

$$\mathcal{Q}_\phi(A) = \mathcal{Q}(A) = E(A).$$

At positive defect the domain has the form (II.185.35); at zero defect it is a coordinate product coset. Constant labels attain equality on every domain in these classifications.

Proof. The chain (II.185.31), the energy envelope, and the strict monotonicity of τ force equality of all three counts. Apply Theorem II.185.15. For a constant label, the two full alternating sums agree on every valid cube pair, so $\mathcal{Q}_\phi(A) = \mathcal{Q}(A)$. $\square$

The last corollary classifies the necessary domain geometry and supplies sufficient constant-label examples. It does not classify all admissible labels. A label classification must retain arbitrary lower-coordinate terms, as the exact endpoint theory in sources 14–15 already emphasizes [202]. Quantitative repair of those labels on the punctured products is a further problem.

II.186 Source 48's chapter: collision-aware derandomization of the restriction step [write]

Added in the ninth write, 7 October 2026.

The next eight sections are source 48's Sections 2–9, on the random restriction of Gowers's Lemma 9.3. *Logically the chapter follows source 17's chapter* (Section II.57), whose Fejér restriction it makes constructive, and source 29's restriction over a fixed field (Section II.119); it is printed here because no number printed earlier may move. Source 48's introduction is Section F.53; its checks, integration notes and questions are in Part V.

What is proved. A rational signed combination F of order- s additive energies on a common set of labels, each label selected independently with a rational probability, is rounded deterministically to a set S with $F(S) \geq \mathbb{E}F$ in $O(C_s n D)$ rational operations and $O((s+1)^2 D)$ storage, $C_s = \binom{s+2}{2}^2$ and D the total size of the ambient groups (Theorem F.53.1). The method is an exact group-algebra inventory of partial energy counts in which a repeated label is selected once (Theorem II.188.2), with a universal Stirling response that turns each probability replacement into translations of coefficient arrays (Theorem II.189.2) and a branch decision by $O(s^2)$ probes (Theorem II.190.1). A distinct-support profile gives the same for a uniformly random subset of any prescribed size, with no factor of that size (Theorem II.191.2); fixed systems of linear equations, with zero coefficients, repeated entries and composite moduli, are handled with column-type compression (Theorem II.192.1). Applied to Lemma 9.3: for $B \subseteq \mathbb{F}_p$, $|B| = \beta p$, and a map ϕ respecting at least αn^{2s-1} balanced tuples, a set B' with respected proportion $1 - \eta$ and $\mathcal{T}_s(B') \geq n^{2s-1}/(2\eta L^{2s})$ is *constructed* with rational arithmetic in $O(C_s n p^5)$ operations, under the threshold of source 17 (Theorem II.194.1); this is source 17's Theorem II.61.2 with its retention, in the unweighted scalar case, made constructive. It answers source 17's Question 9 in its ambient-size form (item II.23). The bounds are polynomial in the ambient size D , not in $\log D$, as source 48 says; no new retention exponent is claimed.

Relation to the formal statements (checked at HEAD 06b30ca52). `lemma_9_3` (Sections08_09.lean, line 77) has the exact companion `lemma_9_3_holds` (Proofs09Restriction.lean, line 376), an existence theorem; the algorithm, its complexity and the constructive theorem are not formalized.

II.187 Sampling counts and the repeated-entry obstruction

Source 48, delivered Section 2.

All groups are finite, abelian, written additively, and supplied in a representation in which translations of a full group-indexed array can be performed in linear time in its size. Products of cyclic groups have this property. Arithmetic complexity counts rational operations separately from the bit cost of those operations.

Definition II.187.1 (Partial energy counts). For nonnegative integers a, b and $h \in K$, let

$$N_{a,b}(S; h) = \# \left\{ (i_1, \dots, i_a, j_1, \dots, j_b) \in S^{a+b} : \sum_{r=1}^a g_{i_r} - \sum_{r=1}^b g_{j_r} = h \right\}.$$

The empty tuple has count $N_{0,0}(S; h) = \mathbf{1}_{h=0}$. Thus $E_s(S; g) = N_{s,s}(S; 0)$.

If S_p is formed by independent label selection, a tuple with distinct-label support A survives

with probability

$$\prod_{i \in A} p_i. \quad (\text{II.187.1})$$

Multiplying p_i once for every occurrence instead computes another quantity. The difference persists even when the underlying group is prime-cyclic and all label images are distinct.

Example II.187.2 (A one-label diagnostic). For one label selected with probability q , its order- s energy is one when selected and zero otherwise. Therefore

$$\mathbb{E}E_s(S_q) = q.$$

The ordinary weighted convolution expression gives q^{2s} . At $q = 1/2$ and $s = 8$ these are $1/2$ and $1/65536$, respectively. An algorithm that substitutes the weighted convolution for the conditional expectation is not merely making a small numerical error.

For a source function $\phi : B \rightarrow H$, use the two labelled images

$$g_x^{\text{dom}} = x \in G, \quad g_x^{\text{gr}} = (x, \phi(x)) \in G \oplus H.$$

Then

$$\mathcal{T}_s(S) = E_s(S; g^{\text{dom}}), \quad \mathcal{R}_s(S, \phi) = E_s(S; g^{\text{gr}}), \quad 0 \leq \mathcal{R}_s(S, \phi) \leq \mathcal{T}_s(S).$$

For $0 < \eta \leq 1$, the useful signed statistic is

$$F_\eta(S) = \mathcal{R}_s(S, \phi) - (1 - \eta)\mathcal{T}_s(S). \quad (\text{II.187.2})$$

Equivalently, every respected tuple contributes η and every unrespected tuple contributes $-(1 - \eta)$.

Lemma II.187.3 (Surplus extraction). *If $F_\eta(S) \geq J > 0$, then*

$$\mathcal{R}_s(S, \phi) \geq (1 - \eta)\mathcal{T}_s(S), \quad \mathcal{T}_s(S) \geq J/\eta.$$

When S is a subset of a group, not a multiset of repeated labels,

$$|S| \geq (J/\eta)^{1/(2s-1)}.$$

Proof. The first assertion is the definition of F_η . Since $\mathcal{R}_s \leq \mathcal{T}_s$, one also has $F_\eta \leq \eta\mathcal{T}_s$. Finally, an additive $2s$ -tuple is determined by its first $2s - 1$ entries: the last entry is forced in the group. Hence $\mathcal{T}_s(S) \leq |S|^{2s-1}$. $\square$

The positivity of a computed surplus is an exact certificate. A negative surplus is only a failed certificate for that sampling distribution; it is not evidence that no useful subset exists.

II.188 An exact group-algebra inventory

Source 48, delivered Section 3.

II.188.1 The finite coefficient algebra

Let $\mathbb{Q}[K]$ have basis symbols $[h]$, $h \in K$, with $[h][h'] = [h + h']$. Work in the finite algebra

$$\mathcal{A}_s(K) = \mathbb{Q}[K][u, v]/(u^{s+1}, v^{s+1}).$$

Use the divided-power basis

$$u^{[a]} = u^a/a!, \quad v^{[b]} = v^b/b!.$$

A coefficient array $A_{a,b}(h)$ represents $\sum A_{a,b}(h)u^{[a]}v^{[b]}[h]$. Divided powers are only a convenient basis over $\mathbb{Q}$; the group may have arbitrary torsion. Their multiplication introduces the integer factors $\binom{a+c}{a}$ and $\binom{b+d}{b}$.

For one label with image g , define the truncated exponential

$$\mathcal{E}_g = \exp(u[g] + v[-g]) = \sum_{0 \leq a, b \leq s} u^{[a]}v^{[b]}[(a-b)g]. \quad (\text{II.188.1})$$

All exponentials and inverses in this section are finite formal operations in a nilpotent ideal. There is no analytic convergence question.

Definition II.188.1 (Bernoulli inventory). Set

$$\mathcal{I}_{\mathbf{p}} = \prod_{i=1}^n (1 + p_i(\mathcal{E}_{g_i} - 1)) = \sum_{a,b,h} A_{a,b}(h)u^{[a]}v^{[b]}[h]. \quad (\text{II.188.2})$$

Theorem II.188.2 (Exact coefficient identity). *For every $0 \leq a, b \leq s$ and $h \in K$,*

$$A_{a,b}(h) = \mathbb{E}N_{a,b}(S_{\mathbf{p}}; h). \quad (\text{II.188.3})$$

In particular, $A_{s,s}(0) = \mathbb{E}E_s(S_{\mathbf{p}}; g)$. The assertion permits repeated entries, repeated group images, probabilities zero or one, and arbitrary finite abelian groups.

Proof. For a deterministic subset S ,

$$\prod_{i \in S} \mathcal{E}_{g_i} = \exp \left(u \sum_{i \in S} [g_i] + v \sum_{i \in S} [-g_i] \right).$$

The coefficient of $u^{[a]}v^{[b]}[h]$ is $N_{a,b}(S; h)$: expansion of the two powers lists the ordered choices of a positive and b negative labels.

Let $X_i = \mathbf{1}_{i \in S_{\mathbf{p}}}$. Independence across labels gives

$$\mathbb{E} \prod_i \mathcal{E}_{g_i}^{X_i} = \prod_i ((1 - p_i) + p_i \mathcal{E}_{g_i}) = \mathcal{I}_{\mathbf{p}}.$$

Taking coefficients proves the identity. Alternatively, a label factor contributes p_i whenever that label occupies a nonempty collection of positions, and contributes one otherwise, which gives (II.187.1) directly. $\square$

II.188.2 Initialization by translations, not convolution

Multiplication of a general inventory by the factor for one label gives

$$A'_{a,b}(h) = A_{a,b}(h) + p_i \sum_{\substack{0 \leq c \leq a, 0 \leq d \leq b \\ c+d > 0}} \binom{a}{c} \binom{b}{d} A_{a-c,b-d}(h - (c-d)g_i). \quad (\text{II.188.4})$$

This is a collection of translations of lower-degree arrays. It is not a convolution of two arbitrary arrays. Start with $A_{0,0}(0) = 1$ and all other coefficients zero, and apply (II.188.4) once per label.

Lemma II.188.3 (Initialization cost). *All coefficients of (II.188.2) can be initialized in $O(C_s n |K|)$ rational operations with $(s+1)^2 |K|$ coefficient locations. The updates can be performed in place, in decreasing order of $a+b$.*

Proof. The number of quadruples (a, b, c, d) with $0 \leq c \leq a \leq s$ and $0 \leq d \leq b \leq s$ is

$$\left(\sum_{a=0}^s (a+1) \right)^2 = \binom{s+2}{2}^2 = C_s.$$

Each nonconstant summand is a weighted translation of an array of length $|K|$. In decreasing total degree, every source row has smaller total degree and has not yet changed. Thus all values on the right side are old values, exactly as required. $\square$

A direct Fourier method can sometimes initialize ordinary weighted energies more quickly. The point here is different: the same small state will continue to give the *correct* conditional expectation after arbitrary indicators have been fixed.

II.189 Local replacement, Stirling coefficients, and bit cost

Source 48, delivered Section 4.

II.189.1 A replacement that never divides by a probability

To replace the probability of one label from p to q , multiply the inventory by

$$R_{p,q}(t) = \frac{1 + q(e^t - 1)}{1 + p(e^t - 1)} = 1 + (q - p)H_p(t), \quad H_p(t) = \frac{e^t - 1}{1 + p(e^t - 1)}, \quad (\text{II.189.1})$$

with $t = u[g] + v[-g]$. The denominator has constant term one even when $p = 0$ or $p = 1$. This removes a common endpoint difficulty in deletion formulas based on division by p or by $1 - p$.

Let $\left\{ \begin{smallmatrix} k \\ j \end{smallmatrix} \right\}$ denote a Stirling number of the second kind, and write

$$H_p(t) = \sum_{k=1}^{2s} h_k(p) t^k / k! \quad \text{modulo terms of degree greater than } 2s.$$

Lemma II.189.1 (Explicit response polynomials). *For $k \geq 1$,*

$$h_k(p) = \sum_{j=1}^k (-p)^{j-1} j! \left\{ \begin{matrix} k \\ j \end{matrix} \right\}. \quad (\text{II.189.2})$$

In particular,

$$h_1 = 1, \quad h_2 = 1 - 2p, \quad h_3 = 1 - 6p + 6p^2, \quad h_4 = 1 - 14p + 36p^2 - 24p^3.$$

The degree is at most $k - 1$, $h_k(0) = 1$, and $h_k(1) = (-1)^{k-1}$.

Proof. Expand the geometric inverse in (II.189.1):

$$H_p(t) = \sum_{j \geq 1} (-p)^{j-1} (e^t - 1)^j.$$

The exponential generating function for ordered partitions into j nonempty blocks is

$$(e^t - 1)^j = j! \sum_{k \geq j} \left\{ \begin{matrix} k \\ j \end{matrix} \right\} \frac{t^k}{k!}.$$

For completeness, its coefficient counts assignments of k labelled positions to j nonempty labelled boxes, which are precisely an unordered set partition followed by an ordering of its blocks. Comparing coefficients proves (II.189.2). The endpoint $p = 1$ follows also from $H_1(t) = 1 - e^{-t}$. $\square$

Theorem II.189.2 (Exact local replacement). *Suppose A is the inventory of the current probabilities. Define $r_0 = 1$ and $r_k = (q - p)h_k(p)$ for $k \geq 1$. Replacing one label image g from probability p to probability q changes the coefficients to*

$$A'_{a,b}(h) = \sum_{c=0}^a \sum_{d=0}^b \binom{a}{c} \binom{b}{d} r_{c+d} A_{a-c,b-d}(h - (c - d)g). \quad (\text{II.189.3})$$

Its cost is $O(C_s|K|)$ rational operations, with an in-place descending implementation. Moreover,

$$A'_{s,s}(0) = A_{s,s}(0) + (q - p)\Delta_g, \quad (\text{II.189.4})$$

$$\Delta_g = \sum_{\substack{0 \leq c, d \leq s \\ c+d > 0}} \binom{s}{c} \binom{s}{d} h_{c+d}(p) A_{s-c,s-d}(-(c - d)g). \quad (\text{II.189.5})$$

Thus the exact slope of the conditional expectation is obtained from $(s + 1)^2 - 1$ coefficient probes.

Proof. The replacement ratio in (II.189.1) cancels exactly the old label factor and inserts the new one. After substituting $t = u[g] + v[-g]$, the coefficient of $u^{[c]}v^{[d]}$ is $r_{c+d}[(c - d)g]$. Multiplication in the divided-power basis gives (II.189.3). The complexity and update order follow from the same count as Lemma II.188.3. Setting $(a, b, h) = (s, s, 0)$ and separating $(c, d) = (0, 0)$ proves the last two identities. $\square$

The response polynomials are universal: neither the group, the label, the other probabilities, nor the source function occurs in them. The only group operation in the formula is translation by an integer multiple of the current image.

II.189.2 Bit complexity is not hidden in an arithmetic-operation count

Proposition II.189.3 (Common-denominator control). *Suppose initially $p_i \in Q^{-1}\mathbb{Z} \cap [0, 1]$ for an integer $Q \geq 1$, and subsequent replacements set probabilities to zero or one. A coefficient of total degree $d = a + b$ has denominator dividing Q^d . The reduced numerators and denominators in the coefficient updates have bit length*

$$O(s(\log(n + 1) + \log Q + \log(s + 1))).$$

A rational objective coefficient adds its own bit length to comparisons. Consequently, for fixed s the arithmetic bounds give polynomial bit-time bounds in the explicit ambient sizes, n , $\log Q$, and the input coefficient bit lengths.

Proof. By Theorem II.188.2, each coefficient is a sum of at most n^d monomials, each a product of probabilities for at most d distinct labels. A common denominator is therefore Q^d .

For an intermediate summand of (II.189.3), the lower coefficient has degree $d - k$, where $k = c + d$ in that formula. The response r_k is a polynomial of degree at most k in grid probabilities. Thus its product with the lower coefficient still has denominator dividing Q^d . The integer coefficients in h_k count ordered set partitions with signs. Their total absolute value is at most $k!2^{k-1}$: order the k elements and insert cuts between consecutive elements, an onto representation of ordered set partitions. Binomial coefficients and the polynomially many summands introduce at most another $O(s \log(s + 1))$ bits. This proves the stated bound.

For arbitrary rational input probabilities, one may use the least common multiple of their denominators as Q . Its logarithm is at most the sum of their denominator bit lengths, so this also gives a polynomial bound. The statement concerns exact reduced arithmetic; it does not say that all these operations have unit bit cost. $\square$

II.190 Deterministic signed-energy rounding

Source 48, delivered Section 5.

The inventory records one expectation at a time, but that expectation is for the *same* random label set in all ambient groups. Thus several inventories can be rounded together without an independence assumption between their statistics.

Theorem II.190.1 (Exact signed rounding). *Let $F(S) = \sum_{\nu=1}^v c_\nu E_s(S; g^{(\nu)})$ with rational c_ν , and let $\mathbf{p} \in ([0, 1] \cap \mathbb{Q})^n$. The following algorithm returns S with $F(S) \geq \mathbb{E}F(S_{\mathbf{p}})$. Initialize the inventory in each group. For each undecided label i , compute its signed derivative*

$$\Delta_i = \sum_{\nu=1}^v c_\nu \sum_{\substack{0 \leq a, b \leq s \\ a+b > 0}} \binom{s}{a} \binom{s}{b} h_{a+b}(p_i) A_{s-a, s-b}^{(\nu)}(-(a-b)g_i^{(\nu)}).$$

Replace p_i by 1 when $\Delta_i \geq 0$ and by 0 otherwise, using the local replacement formula to update each inventory. Return the labels whose final probabilities equal 1.

The cost is $O(C_s n D)$ rational operations and $O((s + 1)^2 D)$ rational storage. The decision itself uses $O((s + 1)^2 v)$ coefficient probes and arithmetic operations after the universal response coefficients have been computed.

Proof. At any stage let J be the current signed expectation and let $p = p_i$. The replacement identity gives the two conditional values

$$J_0 = J - p\Delta_i, \quad J_1 = J + (1 - p)\Delta_i.$$

Their weighted average $(1 - p)J_0 + pJ_1$ equals J . The stated rule chooses an endpoint no smaller than J . All groups use that same endpoint, so the signed combination is genuinely evaluated on one shared subset.

The update identity preserves the exact conditional-expectation invariant. At the last step every probability is zero or one, and the expectation is the deterministic value $F(S)$. Induction therefore proves the inequality. Initial probabilities zero or one are already decided and require no branching.

Initialization and each replacement use at most a constant multiple of $C_s D$ arithmetic operations by the preceding section. Summing over n labels gives the cost. The arrays have $(s + 1)^2 D$ entries. A Stirling table through degree $2s$, the powers of p , and all $h_j(p)$ can be computed in $O(s^2)$ arithmetic operations; this does not increase the stated total. The displayed derivative uses only $(s + 1)^2 - 1$ entries from each group's arrays. $\square$

Corollary II.190.2 (Constructive restriction from a supplied certificate). *Let $B \subseteq G$, $\phi : B \rightarrow H$, and $0 < \eta \leq 1$ be rational. Suppose rational inclusion probabilities satisfy*

$$\mathbb{E}[\mathcal{R}_s(S, \phi) - (1 - \eta)\mathcal{T}_s(S)] \geq J > 0.$$

One can deterministically find $B' \subseteq B$ with

$$\mathcal{R}_s(B', \phi) \geq (1 - \eta)\mathcal{T}_s(B'), \quad \mathcal{T}_s(B') \geq J/\eta,$$

in $O(C_s |B|(|G| + |G||H|))$ rational operations and $O((s + 1)^2(|G| + |G||H|))$ storage.

Proof. Use two inventories, on G and $G \oplus H$, with label images x and $(x, \phi(x))$, and coefficients $-(1 - \eta)$ and 1. Apply the theorem and the surplus-to-restriction implication proved in Section 2. $\square$

Remark II.190.3 (What one signed objective does not imply). A signed combination with a positive lower bound need not preserve each of its summands separately. In particular, adding weighted Fourier or energy statistics does not, by itself, prove several independent lower bounds. The corollary works because one specific surplus implies both quality and retained mass. Other simultaneous conclusions need their own logical implication or a different potential.

II.191 Exact cardinality without a cardinality-sized state

Source 48, delivered Section 6.

Independent rounding does not control the final size. A conventional size-generating variable can introduce a state indexed by every possible cardinality. Here the statistic uses only $2s$ positions, so it is enough to record the number of *distinct undecided labels used by a tuple*. That number never exceeds $2s$, even when the desired subset is very large.

II.191.1 The distinct-support inventory

Suppose I consists of already included labels and U consists of undecided labels, with $|U| = N$. Excluded labels contribute no factor. Put

$$\mathcal{P}_{I,U}(z) = \prod_{i \in I} \mathcal{E}_{g_i} \prod_{i \in U} (1 + z(\mathcal{E}_{g_i} - 1)), \quad \mathcal{E}_g = \exp(u[g] + v[-g]). \quad (\text{II.191.1})$$

The exponentials are the finite truncated polynomials from Section 3. Write $P_{a,b,j}(h)$ for the coefficient of $u^{[a]}v^{[b]}z^j[h]$. It is a nonnegative integer counting ordered partial tuples with difference h and exactly j distinct labels from U ; every remaining label used belongs to I . In particular, $P_{a,b,j} = 0$ for $j > a + b$ and for $j > N$.

For $0 \leq r \leq N$, define a linear functional on these polynomials by

$$\mathcal{L}_{N,r}(z^j) = \begin{cases} \frac{\binom{r}{j}}{\binom{N}{j}} = \frac{(r)_j}{(N)_j}, & 0 \leq j \leq r, \\ 0, & j > r. \end{cases} \quad (\text{II.191.2})$$

The $j = 0$ value is 1, including when $N = r = 0$. No denominator with $j > N$ is used.

Lemma II.191.1 (Exact uniform-subset expectation). *If R is a uniform r -element subset of U , then*

$$\mathbb{E}E_s(I \cup R; g) = \sum_{j=0}^{2s} P_{s,s,j}(0) \mathcal{L}_{N,r}(z^j).$$

The same identity holds for every partial coefficient (a, b, h) .

Proof. A tuple using a particular j -element subset J of U survives exactly when $J \subseteq R$. Its probability is $\binom{N-j}{r-j} / \binom{N}{r} = \binom{r}{j} / \binom{N}{j}$ for $j \leq r$, and zero otherwise. Group the tuples by their distinct undecided support size. Repeated occurrences do not add additional required labels. $\square$

II.191.2 Exclusion and inclusion are polynomial replacements

For a selected undecided label with image g , use the formal shorthand $t = u[g] + v[-g]$. Its old factor is $1 + z(e^t - 1)$. On excluding it, replace that factor by 1; on including it, replace it by e^t . Thus the two multipliers are

$$R_0(t, z) = \frac{1}{1 + z(e^t - 1)} = 1 - zH_z(t), \quad (\text{II.191.3})$$

$$R_1(t, z) = \frac{e^t}{1 + z(e^t - 1)} = 1 + (1 - z)H_z(t). \quad (\text{II.191.4})$$

All inverses are finite nilpotent expansions. Their degree- k divided coefficients are respectively $-zh_k(z)$ and $(1 - z)h_k(z)$, of polynomial degree at most k . The constant coefficient is 1. The translation update from Section 4 therefore applies verbatim, now with polynomial rather than rational entries.

Theorem II.191.2 (Prescribed-size signed rounding). *For every signed energy statistic F and $0 \leq k \leq n$, one can construct a set S of size exactly k with*

$$F(S) \geq \frac{1}{\binom{n}{k}} \sum_{|A|=k} F(A)$$

in $O(C_s(2s+1)^2nD)$ rational operations and $O((s+1)^2(2s+1)D)$ storage. The polynomial coefficient arrays themselves can be maintained over the integers.

Proof. Start with $I = \emptyset$, $U = [n]$, and remaining quota $r = k$. Maintain the inventory (II.191.1) in every group. For one undecided label, the excluded-child expectation J_0 is obtained by the R_0 replacement followed by $\mathcal{L}_{N-1,r}$. The included-child expectation J_1 uses R_1 followed by $\mathcal{L}_{N-1,r-1}$. These are *different evaluators*; applying $\mathcal{L}_{N,r}$ after both updates would give the wrong conditional expectations.

When $0 < r < N$, uniform sampling gives

$$J = (1 - r/N)J_0 + (r/N)J_1.$$

Choose a child of maximum value. If $r = 0$, only exclusion is feasible; if $r = N$, only inclusion is feasible. In those cases the feasible conditional value equals J . The update reduces N by one and reduces r by one exactly when the label is included. It preserves $0 \leq r \leq N$ and never decreases the objective. At termination, $r = N = 0$, exactly k labels have been included, and the conditional value is their statistic.

One may compute each child's top coefficient first, using only lower rows of the current arrays, and then update the whole inventory for the chosen child. No second full child inventory is necessary. Each response and array polynomial has degree at most $2s$. Naive truncated polynomial multiplication takes $O((2s+1)^2)$ rational operations. The translation count is the same $C_s D$ as before. This proves the operation and storage bounds. The response polynomials and initial arrays have integer coefficients; rational arithmetic is needed only for evaluation and comparison of the branch values. $\square$

Remark II.191.3 (Bit bounds and the two sampling laws). An inventory coefficient counts at most n^{a+b} partial tuples. The response-polynomial coefficients are bounded by ordered-partition counts through degree $2s$, so integer intermediate bit lengths are $O(s(\log(n+1) + \log(s+1)))$. The ratios in (II.191.2) have $O(s \log(n+1))$ bits, and there are at most $2s+1$ such terms per evaluation. This gives polynomial bit bounds, including arbitrary rational signed weights.

This theorem starts from a *uniform* prescribed-size subset. It does not assert the same state size for a heterogeneous Bernoulli distribution conditioned on its cardinality. Those distributions have different support-inclusion probabilities; they are not interchangeable.

II.191.3 All sizes and several quotas

With $I = \emptyset$ initially, let c_j be the signed coefficient of z^j in the top zero-difference inventories. Then the exact profile for all uniform subset sizes is

$$\frac{1}{\binom{n}{k}} \sum_{|A|=k} F(A) = \sum_{j=0}^{2s} c_j \frac{\binom{k}{j}}{\binom{n}{j}}, \quad (\text{II.191.5})$$

where the summand is zero when $j > k$ or $j > n$. The initial profile requires only $2s+1$ coefficients, and evaluating it at all k requires $O(sn)$ arithmetic operations using successive binomial-ratio recurrences. Producing a separately rounded set for *every* k can of course require separate rounding runs; this additional output is not free.

More generally, partition the labels into c specified colour classes and prescribe one cardinality in each class. Attach a variable z_a to undecided labels in class a . Every tuple uses at most $2s$ distinct

labels in total, so at most $\binom{2s+c}{c}$ support monomials occur. Uniform sampling with these quotas is independent between colour classes, and the evaluator sends $\prod_a z_a^{j_a}$ to

$$\prod_a \frac{\binom{r_a}{j_a}}{\binom{N_a}{j_a}}.$$

Branch only within the current label's colour, and use the same conditional-average argument. For fixed s, c this gives an exact $O_{s,c}(nD)$ algorithm preserving every prescribed colour quota. This is an exact cardinality constraint, not a general theorem about simultaneously preserving arbitrary spectral statistics.

II.192 General linear systems and column-type compression

Source 48, delivered Section 7.

II.192.1 An inventory for arbitrary position-dependent increments

Let there be t labelled positions. For every input label i and position j , specify an increment $a_{i,j} \in K$. Define

$$N(S) = \#\{(i_1, \dots, i_t) \in S^t : a_{i_1,1} + \dots + a_{i_t,t} = 0\}.$$

For example, a system $A\mathbf{x} = 0$ of d equations over an abelian group G uses $K = G^d$ and $a_{i,j} = A_j x_i$, where A_j is column j . To require a graph map to respect the same system, use $K = (G \oplus H)^d$ and $a_{i,j} = A_j(x_i, \phi(x_i))$.

Work in the squarefree coefficient algebra $\mathbb{Q}[K][z_1, \dots, z_t]/(z_1^2, \dots, z_t^2)$. Put

$$E_i = \prod_{j=1}^t (1 + z_j[a_{i,j}]), \quad I_p = \prod_i (1 + p_i(E_i - 1)).$$

The coefficient of $z_1 \cdots z_t[0]$ is exactly $\mathbb{E}N(S_p)$. Indeed, the term indexed by a subset of positions chooses all their occurrences of one label at once, so its probability is p_i , not a power of p_i .

Theorem II.192.1 (Exact system rounding). *For a fixed number t of positions, a rational signed combination of such system counts on a common label set can be evaluated and rounded without decreasing expectation in $O(3^t nD)$ rational operations and $O(2^t D)$ storage. Here D is the sum of the ambient group sizes of the statistics. No nonzero-coefficient, large-characteristic, or distinct-entry hypothesis is required.*

Proof. Since the variables commute and square to zero, $E_i = \exp(\sum_j z_j[a_{i,j}])$. Write $a_i(U) = \sum_{j \in U} a_{i,j}$ for $U \subseteq [t]$. The same replacement identity with $r_0 = 1$ and $r_k = (q - p_i)h_k(p_i)$ gives

$$A'_S(h) = \sum_{U \subseteq S} r_{|U|} A_{S \setminus U}(h - a_i(U)). \quad (\text{II.192.1})$$

There is no extra factorial: the coefficient of a squarefree monomial z_U in $(\sum_j z_j[a_{i,j}])^{|U|}/|U|!$ is precisely $[a_i(U)]$. Descending subset size gives an in-place update. The number of (S, U) with $U \subseteq S \subseteq [t]$ is 3^t , and there are 2^t coefficient rows. A derivative of the top coefficient requires only $2^t - 1$ probes. Apply the same signed conditional-expectation rule as in Theorem II.190.1. $\square$

The theorem applies over composite moduli because it uses only group addition and formal coefficient identities. A matrix coefficient may be zero or a nonunit. Restrictions involving coefficient invertibility belong to particular analytic existence estimates, not to this exact counting and rounding engine.

II.192.2 Identical column types reduce the exponent

Suppose the positions split into types $1, \dots, c$, with multiplicities $m_1, \dots, m_c$, and $a_{i,j} = b_{i,\tau}$ whenever position j has type τ . Replace the squarefree variables by divided variables $u_\tau^{[a_\tau]}$, $0 \leq a_\tau \leq m_\tau$, and the label exponential by

$$\exp \left(\sum_{\tau=1}^c u_\tau [b_{i,\tau}] \right).$$

A coefficient represents ordered positions within each type. The replacement formula has a product of binomial coefficients, exactly as in the two-type energy inventory. Its operation and storage bounds are

$$O \left(nD \prod_{\tau=1}^c \binom{m_\tau + 2}{2} \right), \quad O \left(D \prod_{\tau=1}^c (m_\tau + 1) \right). \quad (\text{II.192.2})$$

To prove the operation count, sum the number $\prod_\tau (a_\tau + 1)$ of lower multi-indices over all target multi-indices: the sum factors as $\prod_\tau \sum_{a=0}^{m_\tau} (a + 1)$.

A balanced $2s$ -tuple has two types, $+1$ and -1 , both of multiplicity s . Formula (II.192.2) recovers the headline bounds rather than 3^{2s} . In particular, at the source's order $s = 8$, the inventory has 81 rows and the double-index translation budget is 2025, independent of the number of possible sixteen-tuples.

For a system over $\mathbb{F}_p$ of row rank d_0 , row reduction can first replace it by d_0 independent equations, leaving its zero set unchanged. With scalar domain and codomain $\mathbb{F}_p^h$, the graph ambient group then has size $p^{d_0(1+h)}$. This may still be exponentially large in an input dimension. The result is an explicit ambient-size theorem, not a solution to the sparse or high-dimensional representation problem.

II.193 Robust certificates and entirely rational filters

Source 48, delivered Section 8.

II.193.1 A collision-safe perturbation estimate

The phase filters in an analytic existence proof may take irrational values. The inventory algorithm does not need to represent these values exactly: a quantified surplus and a rational approximation are sufficient. The perturbation bound must again count distinct labels, not occurrences as independent selections.

Lemma II.193.1 (Tuple-objective Lipschitz bound). *Let B be a set of n distinct group elements and $m = 2s$. For the objective F_η of (II.187.2), $0 < \eta \leq 1$, and two probability vectors $\mathbf{p}, \mathbf{q} \in [0, 1]^B$,*

$$|\mathcal{J}(\mathbf{p}) - \mathcal{J}(\mathbf{q})| \leq mn^{m-1} \|\mathbf{p} - \mathbf{q}\|_\infty. \quad (\text{II.193.1})$$

More precisely, let I_x be the number of initial additive m -tuples whose distinct-entry support contains x . Then

$$|\mathcal{J}(\mathbf{p}) - \mathcal{J}(\mathbf{q})| \leq \sum_{x \in B} I_x |p_x - q_x|, \quad \sum_x I_x \leq m\mathcal{T}_s(B) \leq mn^{m-1}. \quad (\text{II.193.2})$$

Proof. For numbers $u_1, \dots, u_j, v_1, \dots, v_j \in [0, 1]$, a telescoping product identity gives

$$\left| \prod_{r=1}^j u_r - \prod_{r=1}^j v_r \right| \leq \sum_{r=1}^j |u_r - v_r|.$$

Apply this to the *distinct support* of each tuple. Its signed coefficient in F_η has absolute value at most one. Summing and interchanging the finite sums proves the first part of (II.193.2). Every tuple has at most m distinct entries, and its final entry is determined by the preceding $m-1$, which gives the remaining bounds. $\square$

The refined influence form permits nonuniform precision: high-influence points must be approximated more accurately. The influences themselves are accessible from the derivative formula applied to the unsigned domain energy at the all-one probability vector. This gives an exact, optional precision-allocation certificate rather than an unproved heuristic.

Corollary II.193.2 (Finite-family certificate). *Let Θ be a finite family of probability vectors. Suppose its exact mean surplus is at least J_0 , and rational vectors $\tilde{\mathbf{p}}^\theta$ satisfy $\|\tilde{\mathbf{p}}^\theta - \mathbf{p}^\theta\|_\infty \leq \varepsilon$ for every θ . Searching the rational family and rounding its best member returns a set with*

$$F_\eta(S) \geq J_0 - mn^{m-1}\varepsilon.$$

The cost after constructing the rational family is $O(C_s n D(|\Theta| + 1))$ rational operations, with streaming storage $O((s+1)^2 D + n)$.

Proof. The rational-family average is at least the claimed lower bound by Lemma II.193.1. Its maximum is at least its average. Evaluate every member with the inventory, retain its parameters if best, and apply Theorem II.190.1 to that member. Only the best parameters and its probability vector need be kept. $\square$

II.193.2 A certified rational Fejér table

For $L \geq 2$ define

$$K_L(t) = \frac{1}{L^2} \left| \sum_{j=0}^{L-1} e^{2\pi i j t} \right|^2 = \frac{L + 2 \sum_{j=1}^{L-1} (L-j) \cos(2\pi j t)}{L^2}. \quad (\text{II.193.3})$$

The first expression shows $0 \leq K_L \leq 1$.

Lemma II.193.3 (Rational probability table). *Given integers $p, L \geq 2$ and $Q \geq 1$, an explicitly computable table $q_0, \dots, q_{p-1} \in Q^{-1}\mathbb{Z} \cap [0, 1]$ satisfies*

$$0 \leq K_L(t/p) - q_t < 2/Q \quad (0 \leq t < p). \quad (\text{II.193.4})$$

It can be constructed using exact rational arithmetic and certified finite-series bounds. Its construction is polynomial in $p, L, \log(Q+1)$ in the bit model. No numerical sign oracle is required.

Proof. Put $\varepsilon = 1/(16Q)$. First enclose π in a rational interval of width at most ε . One entirely elementary construction uses

$$\pi = 16 \arctan(1/5) - 4 \arctan(1/239). \quad (\text{II.193.5})$$

For $0 < x < 1$, the alternating partial sums of $\sum_{j \geq 0} (-1)^j x^{2j+1}/(2j+1)$ enclose $\arctan x$, with error bounded by the first omitted term. This follows also by integrating the finite geometric identity for $(1+x^2)^{-1}$. Thus $O(\log(Q+1))$ terms suffice for (II.193.5). The identity itself follows from the tangent addition formula: for $A = \arctan(1/5)$ and $B = \arctan(1/239)$, $\tan(4A) = 120/119$ and $\tan(4A - B) = 1$; the angles locate $4A - B$ in $(0, \pi/2)$.

For each residue r , choose $|r| \leq p/2$. Evaluate the cosine Taylor polynomial at $2r\pi_0/p$, where π_0 is the interval midpoint, until the Lagrange remainder is at most ε . Its argument has absolute value less than 4, so $O(\log(Q+1))$ terms are enough. The uncertainty in the argument contributes at most $\varepsilon/2$, since cosine is 1-Lipschitz. We obtain rational cosine intervals of width at most 3ε .

Use those intervals in the second expression of (II.193.3). The nonconstant weights have total mass $2 \sum_{j=1}^{L-1} (L-j)/L^2 = (L-1)/L < 1$, so the resulting kernel interval has width at most 3ε . Intersect it with $[0, 1]$, obtaining $[\ell_t, u_t]$, and put $q_t = \lfloor Q\ell_t \rfloor / Q$. Then $q_t \leq K_L(t/p)$ and

$$K_L(t/p) - q_t < 3/(16Q) + 1/Q < 2/Q.$$

The value at $t = 0$ may be set exactly to one.

The numbers of series terms are logarithmic in $Q+1$, and their rational numerators, denominators and factorials have polynomial bit lengths in $\log p$ and $\log(Q+1)$. The final table uses $O(pL)$ rational weighted additions after the p cosine intervals are prepared. This proves the claimed bit-time statement. $\square$

The supplied implementation uses this construction. Its numerical cosine comparison in the test suite is supplementary: the certificate comes from the rational interval and the proved remainder bounds.

II.194 A deterministic polynomial-loss restriction theorem

Source 48, delivered Section 9.

II.194.1 Statement and inherited quantitative ingredient

The following theorem is the precise constructive strengthening of the companion's Fejér restriction bound. The source is Gowers's Lemma 9.3 [1]; the polynomial-retention estimate and translated Fejér filter were established in the inspected companion manuscript [207]. We include their short argument so that the algorithmic conclusion does not rest on an unchecked black box.

Theorem II.194.1 (Fully rational constructive restriction). *Let $m = 2s \geq 4$, p be prime, $B \subseteq \mathbb{F}_p$ have size $n = \beta p > 0$, and $\phi : B \rightarrow \mathbb{F}_p$. Let rational $\alpha, \eta \in (0, 1]$ satisfy*

$$\mathcal{R}_s(B, \phi) \geq \alpha n^{m-1}.$$

Define

$$L = \left\lceil \frac{m+1}{\alpha\eta} \right\rceil, \quad d = \alpha + \eta - \alpha\eta, \quad Q = \left\lceil \frac{4mL^m}{d} \right\rceil. \quad (\text{II.194.1})$$

Assume (F.53.3). A deterministic algorithm with only rational probability arithmetic returns $B' \subseteq B$ satisfying (F.53.4)–(F.53.5).

Excluding the separately described rational-kernel precomputation, the complete algorithm uses $O(C_s np^5)$ rational operations and $O((s+1)^2 p^2 + n)$ storage. Once a filter parameter triple is supplied, exact evaluation and rounding use $O(C_s np^2)$ operations. All these operations have polynomial bit cost in the explicit ambient model.

[write] **7 October 2026; sources 47–53.** Against the formal statement (checked at HEAD 06b30ca52). `lemma_9_3` (Sections08_09.lean, line 77) has the exact companion `lemma_9_3_holds` (Proofs09Restriction.lean, line 376), an existence theorem; source 17’s theorem implies its prime case with an explicit threshold (note in source 17’s chapter), and this theorem makes that case constructive with the same constants. The algorithm and its complexity are not formalized.

The complete algorithm first searches p^3 phase triples; the smaller rounding bound is not a bound for that search. For fixed s and $n \leq p$, the complete arithmetic bound is $O_s(p^6)$. It is a substantial reduction from listing the source’s additive sixteen-tuples, but not a practical claim for the very large p required by the worst-case analytic threshold.

II.194.2 Signal of the filter

The Fourier expansion of (II.193.3) is

$$K_L(t) = \frac{1}{L} \sum_{a=-(L-1)}^{L-1} \left(1 - \frac{|a|}{L}\right) e^{2\pi i a t}. \quad (\text{II.194.2})$$

Set

$$S_{m,L} = 1 + 2 \sum_{j=1}^{L-1} (j/L)^m. \quad (\text{II.194.3})$$

Lemma II.194.2 (Signal bounds). For $m \geq 2$,

$$\frac{2L}{m+1} \leq S_{m,L} \leq L.$$

Proof. For the lower bound, the trapezoidal rule for the convex function $f(x) = (1 - x/L)^m$ on $[0, L]$ gives

$$\frac{1}{2}f(0) + \sum_{j=1}^{L-1} f(j) + \frac{1}{2}f(L) \geq \int_0^L f(x) dx = L/(m+1).$$

Twice the left side is $S_{m,L}$. For the upper bound, sum $u^m \leq u$ over the triangular coefficients in (II.194.2). Their sum is L . $\square$

II.194.3 Exceptional tuples and their exact role

Let $\epsilon = (1, \dots, 1, -1, \dots, -1) \in \mathbb{F}_p^m$ with s entries of each sign. An additive tuple $x \in \mathbb{F}_p^m$ is called exceptional when there exists an integer vector $a \in [-L+1, L-1]^m$ such that

$$\sum_j a_j = 0 \pmod{p}, \quad \sum_j a_j x_j = 0, \quad a \bmod p \notin \text{span}(\epsilon).$$

Lemma II.194.3 (Exceptional count). *There are at most*

$$D_0 = (2L - 1)^{m-1} p^{m-2} \quad (\text{II.194.4})$$

exceptional additive tuples. Every nonexceptional tuple has distinct entries. If a vector in the indicated integer box reduces to a multiple of ϵ , it is exactly $t\epsilon$ as an integer vector, with $|t| \leq L - 1$.

Proof. After choosing the first $m - 1$ coefficients, the last is determined modulo p by their sum. The permitted interval contains at most one representative because $p > 2L - 2$. Thus there are at most $(2L - 1)^{m-1}$ coefficient vectors. For a vector outside $\text{span}(\epsilon)$, the equations with rows a and ϵ are independent over $\mathbb{F}_p$, so have p^{m-2} solutions. A union bound proves (II.194.4).

If $x_i = x_j$, the vector $e_i - e_j$ witnesses exceptionality: it is in the coefficient box, its sum is zero, and it cannot be a multiple of ϵ because $m \geq 4$. Finally, if $a \equiv t\epsilon \pmod{p}$, choose $t = a_1$. Each integer $a_j - \epsilon_j t$ is divisible by p and has absolute value at most $2L - 2 < p$, hence vanishes. $\square$

Choose u, v, c uniformly in $\mathbb{F}_p$, and, conditioned on the triple, include $x \in B$ independently with probability

$$p_x^{u,v,c} = K_L((ux + v\phi(x) + c)/p). \quad (\text{II.194.5})$$

The numerator denotes its residue class, so the expression is independent of the chosen lift.

Lemma II.194.4 (Generic survival). *A nonexceptional additive tuple is selected with unconditional probability $L^{-m} S_{m,L}$ if it is respected, and L^{-m} otherwise.*

Proof. The entries are distinct, so conditional independence gives a product of m filter values. Expand them by (II.194.2). Averaging over u, v, c leaves exactly the coefficient vectors satisfying

$$\sum a_j x_j = 0, \quad \sum a_j \phi(x_j) = 0, \quad \sum a_j = 0.$$

By nonexceptionality and Lemma II.194.3, these vectors are $t\epsilon$, $|t| \leq L - 1$. If the tuple is respected, all such t survive, giving $L^{-m} S_{m,L}$. Otherwise only $t = 0$ survives: a nonzero t in that interval cannot annihilate a nonzero scalar discrepancy in $\mathbb{F}_p$. The coefficient for $t = 0$ is L^{-m} . $\square$

II.194.4 A surplus with enough slack for rationalization

Put $T = \mathcal{T}_s(B)$, $R = \mathcal{R}_s(B, \phi)$, and $a = L^{-m}$. Exceptional tuples may be repeated; they are not silently assigned the generic survival probabilities. Let D_R, D_B be the numbers of respected and unrespected exceptional tuples in B . Their sum is at most D_0 . The generic-survival lemma and the bound of one on all survival probabilities give

$$\begin{aligned} \mathbb{E} F_\eta &\geq \eta a S_{m,L} (R - D_R) - (1 - \eta) a (T - R) - (1 - \eta) D_B \\ &\geq a [\eta S_{m,L} R - (1 - \eta) (T - R)] - D_0. \end{aligned} \quad (\text{II.194.6})$$

The last inequality uses $a S_{m,L} \leq 1$.

Let $M = n^{m-1}$. We have $T \leq M$, $R \geq \alpha M$, and $\alpha \eta S_{m,L} \geq 2$ by the choice of L . Consequently,

$$\begin{aligned} a [\eta S_{m,L} R - (1 - \eta) (T - R)] &\geq a [\alpha \eta S_{m,L} - (1 - \eta) (1 - \alpha)] M \\ &\geq (1 + d) a M, \quad d = \alpha + \eta - \alpha \eta > 0. \end{aligned} \quad (\text{II.194.7})$$

The size assumption is exactly a sufficient condition for $D_0 \leq aM/2$. Thus

$$\mathbb{E}_{u,v,c} \mathcal{J}(\mathbf{p}^{u,v,c}) \geq (\tfrac{1}{2} + d)L^{-m}M. \quad (\text{II.194.8})$$

Keeping d instead of discarding it is useful: it pays for a rational filter approximation without weakening the existing retention constant.

II.194.5 Proof of the constructive theorem

Construct the rational kernel table of Lemma II.193.3 with Q from (II.194.1). Replace every probability in (II.194.5) by the corresponding table entry. Uniformly over all triples,

$$\|\tilde{\mathbf{p}}^{u,v,c} - \mathbf{p}^{u,v,c}\|_\infty < 2/Q \leq \frac{d}{2mL^m}.$$

Lemma II.193.1 and (II.194.8) imply

$$\mathbb{E}_{u,v,c} \mathcal{J}(\tilde{\mathbf{p}}^{u,v,c}) \geq (\tfrac{1}{2} + \tfrac{d}{2})L^{-m}M \geq M/(2L^m). \quad (\text{II.194.9})$$

Search all p^3 triples, evaluating their rational objectives with the domain and graph inventories. Keep one attaining the largest value. Theorem II.190.1 produces a deterministic subset with $F_\eta(B') \geq M/(2L^m)$. Lemma II.187.3 proves (F.53.4) and the first bound of (F.53.5). Since $L \leq (m+2)/(\alpha\eta)$, its second bound follows.

For one triple, the total group size is $p+p^2$ and initialization costs $O(C_s n(p+p^2))$. Multiplication by p^3 and one final rounding gives $O(C_s np^5)$. Streaming gives the stated storage. All rational probabilities have common denominator Q , so Proposition II.189.3 applies. Kernel precomputation is polynomial by Lemma II.193.3; under the hypothesis $L < (p+2)/2$, its $O(pL)$ table stage is also polynomial in the ambient size. This completes the proof of Theorem II.194.1.

II.194.6 An explicit search–threshold tradeoff

Corollary II.194.5 (Two phase parameters). *The same retention conclusion and rational denominator Q are obtained by searching only $(u, v) \in \mathbb{F}_p^2$, setting $c = 0$, provided the size hypothesis is replaced by*

$$p > 2L - 2, \quad p \geq 2L^m(2L - 1)^m \beta^{-(m-1)}. \quad (\text{II.194.10})$$

The arithmetic cost is then $O(C_s np^4)$.

Proof. Define exceptionality without the coefficient-sum condition. There are at most $(2L - 1)^m$ possible coefficient vectors; the same rank-two count gives exceptional mass at most $(2L - 1)^m p^{m-2}$. Nonexceptional tuples still have distinct entries and the same principal-line classification. Averaging over u, v now imposes only the domain and image equations, and the generic survival probabilities are unchanged. The larger threshold makes the new exceptional bound at most $M/(2L^m)$. Repeat (II.194.7)–(II.194.9). There are now p^2 parameter choices instead of p^3 . $\square$

This tradeoff is not simultaneously a better density threshold and a faster search: it saves one factor p in search at the cost of one factor $2L - 1$ in the sufficient ambient-size threshold.

II.194.7 Exact data can outperform the worst-case certificate

For a concrete rational kernel table, the algorithm computes

$$\max_{u,v,c} \mathcal{J}(\tilde{\mathbf{p}}^{u,v,c}) \quad \text{and, optionally,} \quad p^{-3} \sum_{u,v,c} \mathcal{J}(\tilde{\mathbf{p}}^{u,v,c})$$

exactly, including every degenerate tuple. A positive maximum is sufficient even when (F.53.3) fails spectacularly. The threshold is a uniform sufficient condition, not a prerequisite for running the algorithm. No separate enumeration of the exceptional set is required for this exact-data certificate.

The same exact evaluation supports finite searches over L , η , or a user-specified rational kernel family, as long as their costs are counted. This implements a finite optimization procedure; it is not a proof that Fejér kernels or the chosen grid are globally optimal filters.

II.195 Sources 53, 51 and 52: Section 13 square geometry and Stage 13.7 [write]

Added in the ninth write, 7 October 2026.

The next twenty-four sections are source 53's Sections 2–8, source 51's Sections 2–10 and source 52's Sections 2–9. They open this report's first chapter on Gowers's Section 13, the extraction of a bilinear square from the eight-arrangement data (Stages 13.4–13.9 and Corollary 13.10), *as encoded in the repository*: all three study the Lean statements, not the paper's estimates. They are Report279–281 of the series of sources 26, 32, 41 and 42 and do not cite one another. Their introductions are Sections F.58, F.56 and F.57; their questions, verification notes and source maps are in Part V.

What source 53 proves (the spine). For $0 < a_0 \leq 1$ put $N_0(a_0) = \lceil (8/c(a_0))^{1/\gamma(a_0)} \rceil$ with $c = \zeta_{13}/2$ and $\gamma(a) = 2^{-386} a^{1856} 2^{-13Q(a)}$, the formal **section13SquareExponent**. For every prime $N \geq N_0(a_0)$ and every Section 13 context of actual density $a \geq a_0$ there are proper progressions V, W of the same nonzero step and length $k \geq D(a)N^{4\gamma(a)} - 1 \geq N^{3\gamma(a)}$, $D(a) > 1/2$, and $B \subseteq A \cap (V \times W)$ of mass at least $2^{-137} a^{704} k^2$ on which the original phase is bilinear (Theorem F.58.1). Seven sufficient budgets collapse exactly to the seventh (Proposition II.199.1); the short Stage 13.6 witness ($\ell R \in \{m, m+1\}$) is recovered from the construction itself, not from its exported type (Section II.200); rectification inside a short parent, balanced aligned parents and an endpoint-contained cover give the side $\lfloor L/2 \rfloor^g - 1$ (Proposition II.201.4). The analytic inputs are imported, explicitly, as proved Lean theorems at the pin. The threshold is sufficient, not optimal.

What source 51 proves. For every large prime N there are a genuine context of varying density $\alpha = \lfloor N/8 \rfloor / N$, complete Stage 13.4–13.5 data, all six budgets of the budget-qualified extraction theorem and Stage 13.6–13.9 data with $\ell R \leq \ell Q$, whose retained set $S \times U$ has $(S - S) \cap (U - U) = \{0\}$: no common-step square of side two, so the grid predicate **Corollary1310Grid** fails for these witnesses (Theorem F.56.1). This refutes a universal implication from retained data, not existential extraction or **corollary_13_10**, as source 51 says. Its repair by reselection (Theorem II.211.3, Corollary II.211.4) gives exponent $2\gamma(\alpha)$; source 53's Proposition II.201.4 is the same argument without the last weakening.

What source 52 proves. The two all-moduli encodings preserved in the repository, **lemma_13_7_without_domain** and **lemma_13_7_without_prime_assumption**, are false: over $N = 2^k$, $k = 939 \cdot 2^{1892} + 5$, the even residues $T = \{2j : j < N/16\}$ carry the order-eight Freiman map

$2j \mapsto j$, which every ambient affine map meets in at most one point, so the Stage 13.7 length and mass requirements are incompatible (Theorem F.57.1); variants add genuine Stage 13.4 data, unit input steps, and (with a non-unit step) all six numerical budgets with $R \subseteq Q$ (Theorems F.57.2, II.217.1); the mechanism persists over every fixed prime-power family (Theorem II.218.1). The prime-corrected `lemma_13_7` is untouched, and the combination of a unit step with the full package is left open, as source 52 says.

Relation to the formal statements (checked at HEAD 06b30ca52). Since the pins the repository has repaired its Section 13 chain: `corollary_13_10` (`Sections12_13.lean`, line 686) is an exact companion, `corollary_13_10_holds` (`Proofs13ConstructedChain.lean`, line 31), consuming construction geometry exported by the preceding lemmas; the older interface is preserved as `corollary_13_10_without_construction_geometry` (`Sections12_13.lean`, line 669), unasserted; `theorem_13_12` is exact (`theorem_13_12_holds`, `Proofs13EndpointPrintedTheorem.lean`, line 73); and the density-uniform extraction has an explicit threshold, `section13_complete_square_extraction_explicit` (`Proofs13ExplicitGeometricThreshold.lean`, line 24), with side $N^{\gamma(\delta)}$ at the lower density. Source 53's closed threshold, actual-density side and mass, source 51's two theorems and source 52's negations are not formalized (notes in Sections F.58, F.56 and F.57).

II.196 The context and the quantifiers

Source 53, delivered Section 2.

A progression $P = P(x, d; L)$ in $\mathbb{F}_N$ is the indexed sequence $x + jd$ for $0 \leq j < L$. Its carrier is denoted by P when no confusion is possible, and its length is $\ell P = L$. It is proper if the indexing is injective; then $|P| = \ell P$. A common-step square means the product of two proper equal-length progressions whose steps agree and are nonzero. Their starting points need not agree.

A map is order-eight Freiman on a set if it preserves equality between sums of eight elements, with repetitions allowed. A function $\phi : \mathbb{F}_N^2 \rightarrow \mathbb{F}_N$ is separately order-eight Freiman on A when each horizontal and vertical restriction has this property on the corresponding section of A .

An eight-arrangement at height h consists of sixteen horizontal coordinates x_i and sixteen vertical coordinates y_i such that

$$\sum_{i=1}^8 x_i = \sum_{i=9}^{16} x_i, \quad (x_i, y_i), (x_i, y_i + h) \in A \quad (1 \leq i \leq 16).$$

It is respected by ϕ if the same equality of sums holds for the sixteen values $\phi(x_i, y_i + h) - \phi(x_i, y_i)$. Let $C(h)$ count all these arrangements, and let $C_\phi(h)$ count the respected ones. These are ordinary finite counts, without normalization.

Definition II.196.1 (Section 13 context). The context consists of $A \subseteq \mathbb{F}_N^2$, a function $\phi : \mathbb{F}_N^2 \rightarrow \mathbb{F}_N$, and $a \in (0, 1]$, with

$$|A| = aN^2, \quad (1 - 2^{-44}) \sum_h C(h) \leq \sum_h C_\phi(h), \quad (\text{II.196.1})$$

and the separate order-eight Freiman property. These are exactly the hypotheses of the source record `Section13Context`; the auxiliary parameter η is fixed there to 2^{-44} .

The actual density is $a = |A|/N^2$ and may vary with N . Bilinearity in the conclusion means (F.58.9) on B ; it does not assert equality with ϕ everywhere on $V \times W$.

II.196.1 Why the uniform order of quantifiers matters

The older fixed-exact-density statement has the form: for every fixed $a > 0$, there is $N_0(a)$ such that the conclusion holds whenever $|A| = aN^2$ and $N \geq N_0(a)$ is prime. Taken only as an exported statement, it can be eventually vacuous for each fixed $0 < a < 1$. If a is irrational, the cardinality identity is impossible. If $a = r/s$ in lowest terms with $0 < r < s$, it forces $s \mid N^2$; for prime N , this requires $s = N$ or $s = N^2$, and hence permits at most one prime. This observation neither disproves the fixed-density theorem nor invalidates its constructive proof.

The newer uniform source theorem and Theorem F.58.1 have the nonvacuous order

$$\forall a_0 \in (0, 1] \exists N_0(a_0) \forall N \text{ prime} \forall (A, \phi, a), \quad a_0 \leq a \leq 1, \quad N \geq N_0(a_0) \implies \text{conclusion}.$$

The old quantifier issue is therefore not a remaining defect in the pinned uniform source theorem. No threshold here is obtained by substituting the varying density into a pointwise existential statement.

II.197 Scalar parameters and imported analytic inputs

Source 53, delivered Section 3.

The Section 10 cutoff and the Stage 13.4 scalar bounds are

$$\theta(a) = 2^{-59} a^{176}, \tag{II.197.1}$$

$$t(a) = 2^{-1882} \theta(a)^{10477} = 2^{-620025} a^{1843952}, \tag{II.197.2}$$

$$b(a) = 2^{1882} \theta(a)^{-10479} = 2^{620143} a^{-1844304}. \tag{II.197.3}$$

Also put

$$W(a) = 2^{135} a^{-704}, \quad \kappa(a) = \lceil K(a) \rceil, \quad z(a) = \frac{2^{-155\kappa(a)} (a^{32}/16)^{18\kappa(a)}}{\kappa(a)}. \tag{II.197.4}$$

Thus z includes both the source's natural ceiling and division by that integer. In particular z is not being replaced by an unrounded radius.

II.197.1 Monotonicity and the relative exponent margin

On any interval $[a_0, 1]$, the functions $\theta, t, c, \beta, f, g, z, \gamma$ are nondecreasing, while b, W, K, Q are nonincreasing. The explicit powers prove most assertions immediately. For c , increasing a decreases K , increases 2^{-228K} , and increases a^{576K} : the base increases while the positive exponent decreases and the base is at most one. For z , κ decreases, so all three factors $2^{-155\kappa}$, $(a^{32}/16)^{18\kappa}$ and $1/\kappa$ increase. This reasoning does not require continuity of a ceiling-dependent function. The pinned parameter files supply these same monotonicity results.

We shall use

$$0 < c(a) \leq \frac{1}{2}, \quad 0 < 2\beta(a) \leq 1, \quad 0 < f(a), g(a) \leq 1. \tag{II.197.5}$$

For an actual positive integer spectrum count $j \leq b(a)$, write

$$u(a, j) = \frac{t(a)^2}{16j}, \quad v(j) = 2^{-12j}, \quad w(j) = 2^{-11j}. \tag{II.197.6}$$

The letter j is the actual integer count. It is distinct from the much larger real function $Q(a)$ in (F.58.1).

Lemma II.197.1 (Uniform relative margin). *For every $0 < a \leq 1$ and $1 \leq j \leq b(a)$,*

$$2\beta(a) < u(a, j)v(j) \leq u(a, j)w(j). \quad (\text{II.197.7})$$

Proof. The explicit expressions give $b(a) < Q(a)/2$ and $t(a)b(a) = \theta(a)^{-2} \geq 1$. Thus, putting $Q = Q(a)$,

$$u \geq \frac{1}{16b(a)^3} \geq \frac{1}{2Q^3}, \quad v \geq 2^{-6Q}.$$

For $Q \geq 2$, the elementary inequality $4Q^3 < 2^{7Q}$ implies $2 \cdot 2^{-13Q} < 2^{-6Q}/(2Q^3) \leq uv$. For example $Q \leq 2^Q$ for $Q \geq 1$ gives $4Q^3 \leq 2^{2+3Q} < 2^{7Q}$ for $Q \geq 2$. Finally $v \leq w$. This is also the source's proved factor-two recurrence margin. $\square$

II.197.2 The analytic interface used in the proof

We import the following source results. They are not reproved by the companion diagnostics. The source files and declarations are identified in Appendix V.204.

Initial progression and recurrence. For a context and the admissible cutoff $\theta(a)$, Stage 13.4 constructs a positive integer $j \leq b(a)$ and a proper nonzero-step progression P . With

$$x = \frac{t(a)}{64\pi} N^{u(a,j)}, \quad p = \lfloor x \rfloor,$$

one has $\ell P = p$ or $\ell P = p - 1$, together with the full spectral and weighted-height conditions of Stage 13.4. For an odd prime N , Stage 13.5 then constructs the genuine progression $Q_{\text{st}} \subseteq P$ and its spectral recurrence data if

$$T_2(j) < \ell P, \quad 8 \leq (\ell P)^{v(j)}, \quad 160 + 2a^{32} \leq a^{32}(\ell P)^{v(j)}. \quad (\text{II.197.8})$$

The exact simultaneous quadratic threshold is

$$T_2(j) = 2^{(1+2^{321})2^{11(j-1)}} \quad (j \geq 1). \quad (\text{II.197.9})$$

It is an integer and is increasing in j . The Stage 13.5 output is proper with nonzero step, has $\ell Q_{\text{st}} \geq (\ell P)^{v(j)}/2$, and has the required recurrence error $(\ell P)^{-w(j)}$. These assertions summarize the numerical interface, not replacements for the other stage conditions; the imported theorem supplies all of those conditions.

All-density Stage 13.6. For those genuine stages, and an integer m satisfying

$$\begin{aligned} 1 \leq m, \quad m^2 \leq N, \quad m+1 \leq \ell Q_{\text{st}}, \quad c(a)N^{\beta(a)} \leq m, \\ (\ell P)^{-w(j)} \leq z(a)/m, \end{aligned} \quad (\text{II.197.10})$$

the all-density constructor produces genuine Stage 13.6 data with a proper progression R , nonzero step $R = \text{step } Q_{\text{st}}$, and

$$c(a)N^{\beta(a)} \leq \ell R \leq \ell Q_{\text{st}}. \quad (\text{II.197.11})$$

There is no hypothesis $a \leq 1/6$. Moreover its actual construction may retain

$$\ell R = m \quad \text{or} \quad \ell R = m + 1. \quad (\text{II.197.12})$$

This last assertion is a strengthened existential selection, verified from the same witness construction in Section II.200. It is not inferred for every witness satisfying only (II.197.11).

Rows and coefficients. Stage 13.7 constructs a proper nonzero-step progression $S \subseteq R$ of length

$$L = \ell S \geq (\ell R)^{f(a)}. \quad (\text{II.197.13})$$

If $L \geq W(a)$, Stage 13.8 supplies $C \subseteq A$, a set of heights $J_0 \subseteq Q_{\text{st}}$, an offset y , and scalar order-eight Freiman maps A_c, B_c on J_0 such that, with $\delta = 2^{-135}a^{704}$,

$$\begin{aligned} C &\subseteq S \times (J_0 + y), & |C| &\geq \delta L \ell Q_{\text{st}}, \\ \phi(x, z) &= A_c(z - y) + B_c(z - y)x & ((x, z) \in C). \end{aligned} \quad (\text{II.197.14})$$

The scalar Freiman assertions follow by projecting the coefficient pair in the exact Stage 13.8 predicate. These results are valid on the full density range.

The further all-scale coefficient-localization theorem will be stated precisely in Section II.201. It is the only analytic ingredient within the geometric argument there.

II.198 Seven sufficient quantitative budgets

Source 53, delivered Section 4.

Fix $a_0 \in (0, 1]$. A subscript zero on a scalar function means evaluation at a_0 . Define the fixed quantities

$$q = \lfloor b_0 \rfloor, \quad d_0 = \frac{t_0}{64\pi}, \quad d = d_0/2, \quad u_0 = \frac{t_0^2}{16q}, \quad v_0 = 2^{-12q}. \quad (\text{II.198.1})$$

Here $q \geq 1$, $0 < d < 1$, and $u_0, v_0 > 0$. Every actual integer count $j \leq b(a)$ satisfies $j \leq q$, so $u(a, j) \geq u_0$ and $v(j) \geq v_0$.

Put

$$\begin{aligned} H &= \max\{8, \lceil 160a_0^{-32} + 2 \rceil\}, \\ L_* &= \max\{T_2(q) + 1, \lceil H^{1/v_0} \rceil\}, \\ X_* &= \max\{8, W_0, 4 \cdot 8^{1/g_0}\}. \end{aligned} \quad (\text{II.198.2})$$

Both H and L_* are natural numbers. The separate $+1$ in $T_2(q) + 1$ preserves the strict recurrence requirement.

Consider seven positive real thresholds:

$$\begin{aligned} R_1 &= 3, & R_2 &= ((L_* + 2)/d_0)^{1/u_0}, \\ R_3 &= (1/c_0)^{1/\beta_0}, & R_4 &= (4/d)^{1/\beta_0}, \\ R_5 &= (1/(z_0 d))^{1/\beta_0}, & R_6 &= (X_*^{1/f_0}/c_0)^{1/\beta_0}, \\ R_7 &= (8/c_0)^{1/\gamma_0}. \end{aligned} \quad (\text{II.198.3})$$

Their dependence is solely on a_0 . We first explain all seven budgets, before proving in Section II.199 that R_7 strictly dominates the others.

Proposition II.198.1 (Constructive consequences of the seven budgets). *If N is prime and $N \geq \max_i R_i$, then every context with $a_0 \leq a \leq 1$ admits genuine Stages 13.4–13.8 satisfying (II.197.14). They may be selected with the short-parent alternative (II.197.12), and they satisfy*

$$2(\ell R - 1) < N, \quad L \geq (\ell R)^{f(a)} \geq (c(a)N^{\beta(a)})^{f(a)} \geq X_* \geq W(a). \quad (\text{II.198.4})$$

Proof. Use the imported Stage 13.4 constructor, and abbreviate $u = u(a, j)$, $v = v(j)$, $w = w(j)$. Since $N \geq R_1$, the prime is odd. Since $N \geq R_2$,

$$x = \frac{t(a)}{64\pi} N^u \geq d_0 N^{u_0} \geq L_* + 2.$$

The floor and optional predecessor each lose at most one unit, so $\ell P > x - 2 \geq L_*$; in particular $\ell P \geq L_*$. Thus

$$T_2(j) \leq T_2(q) < \ell P, \quad (\ell P)^v \geq (\ell P)^{v_0} \geq H.$$

Because $H \geq 160a_0^{-32} + 2 \geq 160a^{-32} + 2$, all three conditions in (II.197.8) follow. Apply the Stage 13.5 constructor.

For later use the same budget gives the stronger fixed-coefficient statement

$$d_0 N^u \geq d_0 N^{u_0} \geq L_* + 2 \geq 4.$$

Hence the two units of rounding loss imply

$$\ell P \geq (d_0/2) N^u = d N^u. \quad (\text{II.198.5})$$

It matters here that the smaller fixed coefficient d_0 itself has the lower bound four, not merely the actual pre-floor quantity x .

Set $A_* = c(a) N^{\beta(a)}$ and $m = \lceil A_* \rceil$. The third threshold gives $A_* \geq c_0 N^{\beta_0} \geq 1$. Since $c(a) \leq 1/2$,

$$1 \leq m \leq A_* + 1 \leq 2A_* \leq N^{\beta(a)}, \quad m^2 \leq N^{2\beta(a)} \leq N. \quad (\text{II.198.6})$$

For either $s = v$ or $s = w$, (II.198.5), $0 < d < 1$, $s \leq 1$, and Lemma II.197.1 give

$$(\ell P)^s \geq d^s N^{us} \geq d N^{2\beta(a)}.$$

The fourth threshold implies $d N^{\beta(a)} \geq 4$, whence

$$\ell Q_{\text{st}} \geq \frac{1}{2} (\ell P)^v \geq 2 N^{\beta(a)} \geq m + 1.$$

The fifth implies $z(a) d N^{\beta(a)} \geq 1$, whence $z(a) (\ell P)^w \geq N^{\beta(a)} \geq m$, equivalently the last condition in (II.197.10). Every integer and Bohr budget is now verified. Use the all-density constructor with this m , retaining its actual length alternative (II.197.12).

The sixth threshold yields

$$A_* \geq c_0 N^{\beta_0} \geq X_*^{1/f_0} \geq 1.$$

Since $f(a) \geq f_0$, it follows that $A_*^{f(a)} \geq X_*$. Also $A_* \geq 8$, because $X_* \geq 8$ and $f_0 \leq 1$. Thus $m \geq 8$, and

$$2(\ell R - 1) \leq 2m < m^2 \leq N.$$

Apply Stages 13.7 and 13.8. Their hypotheses hold because $L \geq (\ell R)^{f(a)} \geq A_*^{f(a)} \geq X_* \geq W_0 \geq W(a)$. Their output gives (II.197.14) and (II.198.4). This proves the proposition; the seventh budget will pay the final clean-exponent absorption. $\square$

II.198.1 The same budgets also verify the existing square API

The sixth budget is deliberately compatible with the previously exported indexed-span square theorem. Put $X = (\ell R)^{f(a)} \geq X_*$. Its three scale hypotheses are

$$8 \leq X, \quad W(a) \leq X, \quad 8 \leq \left(\left\lfloor \frac{\lfloor X \rfloor}{2} \right\rfloor \right)^{g(a)}.$$

The inner expression is natural integer division. For $X \geq 8$, it is at least $X/4$, so the last inequality follows from $X_* \geq 4 \cdot 8^{1/g_0}$ and $g(a) \geq g_0$. The source theorem then gives a square with actual-density mass and side at least

$$\left(\frac{(c(a)N^{\beta(a)})^{f(a)}}{4} \right)^{g(a)/2} - 1 \geq \frac{c_0}{4} N^{2\gamma(a)} - 1. \quad (\text{II.198.7})$$

Indeed $(c(a)^{f(a)}/4)^{g(a)/2} \geq c(a)/4 \geq c_0/4$. From $N \geq R_7$ and monotonicity of γ , $(c_0/4)N^{\gamma(a)} \geq 2$, so (II.198.7) is at least $N^{\gamma(a)}$.

This gives a direct actual-density refinement using only the already exported square API. The stronger theorem uses the same budgets but retains more of the selected geometry. No square-root loss from (II.198.7) will be imposed on the sharper route.

II.199 The seventh threshold dominates exactly

Source 53, delivered Section 5.

We now prove that the apparent maximum in (II.198.3) is exactly R_7 , on the entire interval $0 < a_0 \leq 1$.

Proposition II.199.1 (Exact threshold collapse). *For every $0 < a_0 \leq 1$, $R_i < R_7$ for $1 \leq i \leq 6$. Consequently*

$$\left\lceil \max_{1 \leq i \leq 7} R_i \right\rceil = \lceil R_7 \rceil = N_0(a_0).$$

Proof. Throughout this proof only, suppress the subscript zero, set $h = \log_2(1/a_0) \geq 0$, and write $K, Q, b, t, c, \beta, f, g$ for their values at a_0 . Then

$$\begin{aligned} K &= 2^{114+320h}, & Q &= 2^{1048576+2097152h}, \\ b &= 2^{620143+1844304h}, & t &= 2^{-620025-1843952h}, \\ f &= 2^{-100-448h}, & g &= 2^{-284-1408h}. \end{aligned}$$

Introduce logarithmic losses

$$C = \log_2(1/c) = 1 + K(228 + 576h), \quad D_0 = \log_2(1/d_0), \quad D_1 = \log_2(1/d) = D_0 + 1. \quad (\text{II.199.1})$$

The logarithm of the seventh term is exactly

$$E = \log_2 R_7 = \frac{4(C+3)}{\beta f g} = (C+3)2^{13Q+386+1856h}. \quad (\text{II.199.2})$$

All comparisons below concern positive quantities.

We use the elementary estimate

$$1 + h \leq 2^{2h} \quad (h \geq 0), \quad (\text{II.199.3})$$

which follows from $\log 2 \geq 1/2$ and $e^x \geq 1 + x$. Since $\pi < 4$,

$$\begin{aligned} D_0 &< 620033 + 1843952h, \\ D_1 &< 620034 + 1843952h, \\ D_1 + 2 &< 620036 + 1843952h < 2^{21}(1 + h) < C. \end{aligned} \quad (\text{II.199.4})$$

The last inequality follows from $C > 2^{114}(228 + 576h)$ by comparing both positive coefficients.

Terms one, three and four. Equation (II.199.2) gives $E > 4 > \log_2 3$. The third logarithm is C/β , and

$$\frac{E}{C/\beta} = \frac{4(C + 3)}{fgC} > 4.$$

The fourth logarithm is $(D_1 + 2)/\beta < C/\beta$. Thus all three terms are strictly below R_7 .

Term five and the rounded radius. The exact radius parameter is $\kappa = \lceil K \rceil$. Because $K \geq 1$, $1 \leq \kappa \leq K + 1 \leq 2K$. From (II.197.4),

$$Z := \log_2(1/z) = \kappa(227 + 576h) + \log_2 \kappa.$$

For a positive integer κ , $\kappa \leq 2^\kappa$, so $\log_2 \kappa \leq \kappa$. Therefore

$$Z \leq \kappa(228 + 576h) \leq 2K(228 + 576h) = 2(C - 1) < 2C. \quad (\text{II.199.5})$$

The fifth logarithm is $(Z + D_1)/\beta < 3C/\beta < E$. This pays the integer ceiling and the denominator in the source radius without treating either as negligible.

Term six and the exact branch of the maximum. Since $2^{1408h} \geq 1 + 704h$,

$$1/g = 2^{284} 2^{1408h} \geq 2^{284}(1 + 704h) > 135 + 704h = \log_2 W.$$

Also $3 = \log_2 8 < 2 + 3/g$. Consequently

$$X_* = 4 \cdot 8^{1/g}, \quad \log_2 X_* = 2 + 3/g. \quad (\text{II.199.6})$$

Writing $E_6 = \log_2 R_6$ and multiplying by the positive common denominator gives

$$E_6 \beta f g = C f g + 2g + 3 \leq C + 5 < 4(C + 3) = E \beta f g.$$

Every branch of the maximum defining X_* is therefore accounted for.

Term two and the strict recurrence buffer. Here $q = \lfloor b \rfloor \geq 512$. Since $160 \cdot 2^{32h} + 2 \geq 162$,

$$H = \lceil 160 \cdot 2^{32h} + 2 \rceil \leq 160 \cdot 2^{32h} + 3 \leq 163 \cdot 2^{32h} < 2^{8+32h}. \quad (\text{II.199.7})$$

Let $T = T_2(q)$ and $J = \log_2 T$. Retaining the exact factor $1 + 2^{321}$,

$$J = (1 + 2^{321})2^{11(q-1)} < 2^{11q+311} \leq 2^{12q}. \quad (\text{II.199.8})$$

Put $M = \max\{T, H^{2^{12q}}\}$. Each entry is greater than one, and both rounding losses in L_* give

$$T + 1 \leq M + 1, \quad \lceil H^{2^{12q}} \rceil \leq M + 1, \quad L_* + 2 \leq M + 3 < 4M.$$

In particular the strict source requirement has not been replaced by a non-strict one. Taking logarithms and using (II.199.7)–(II.199.8),

$$\log_2(L_* + 2) < 2 + 2^{12q}(8 + 32h).$$

Then (II.199.4) and (II.199.3) yield

$$\begin{aligned} \log_2(L_* + 2) + D_0 &< 2^{12q}(10 + 32h + D_0) \\ &< 2^{12q}(620043 + 1843984h) \\ &< 2^{12q+21}(1 + h) \leq 2^{12q+21+2h}. \end{aligned} \quad (\text{II.199.9})$$

The reciprocal exponent is exactly $1/u_0 = q 2^{1240054+3687904h}$. Hence $E_2 = \log_2 R_2$ satisfies

$$\begin{aligned} E_2 &< q 2^{12q+1240075+3687906h} \\ &\leq 2^{12q+1860218+5532210h}, \end{aligned} \quad (\text{II.199.10})$$

using $q \leq b = 2^{620143+1844304h}$. Now

$$\frac{b}{Q/2} = 2^{-428432-252848h} < 1,$$

so $12q < 6Q$. The remaining affine expression is bounded by

$$1860218 + 5532210h < 2^{23}(1 + h) \leq 2^{23+2h} < Q. \quad (\text{II.199.11})$$

Together, (II.199.10)–(II.199.11) imply

$$E_2 < 2^{7Q} < (C + 3)2^{13Q+386+1856h} = E.$$

This completes all six strict comparisons, including $h = 0$. □

[write] 7 October 2026; sources 47–53. Checked at the write in exact integer arithmetic: $\theta = 2^{-59}a^{176}$, $t = 2^{-1882}\theta^{10477} = 2^{-620025}a^{1843952}$ and $b = 2^{1882}\theta^{-10479} = 2^{620143}a^{-1844304}$; $\gamma = \beta fg/4 = 2^{-386}a^{1856}2^{-13Q}$; $b/(Q/2) = 2^{-428432-252848h}$; $1/u_0 = q 2^{1240054+3687904h}$; $D_1 + 2 < 620036 + 1843952h < 2^{21}(1 + h)$, $1860218 + 5532210h < 2^{23}(1 + h)$ and $A_1 \leq 2^{124+322h}$ for $h = 0, \dots, 200$ and as linear forms; $2^{1408h} \geq 1 + 704h$ and $1 + h \leq 2^{2h}$ on the same range. These finite checks do not replace the proof.

Only the final natural ceiling is now needed: $N \geq \lceil R_7 \rceil$ implies $N \geq R_7$ and every previous budget. There is no additional unit at this stage; the only strict recurrence condition was already protected by $T_2(q) + 1$.

Proof of the threshold comparison in Corollary F.58.2. With the notation above, let $A_1 = C + 3 = 4 + K(228 + 576h)$. Since $K \geq 1$,

$$A_1 \leq K(232 + 576h) < K 2^{10}(1 + h) \leq 2^{124+322h}.$$

Equation (II.199.2) therefore gives

$$\log_2 \log_2 R_7 < 13Q + 510 + 2178h < 13Q + 2^{12}(1 + h) < 14Q.$$

For the last step use $2^{12}(1 + h) \leq 2^{12+2h} < Q$. Exponentiating twice gives $R_7 < 2^{2^{14}Q}$. Monotonicity of the natural ceiling proves $N_0 \leq N_{\text{safe}}$; no strict comparison after rounding is required. $\square$

II.200 Why the all-density witness is short

Source 53, delivered Section 6.

The strength of the geometric argument depends on preserving information in an actual selection. We verify that this is available throughout $0 < a \leq 1$.

The all-density initial-stage wrapper in [Proofs13AllDensityExtraction.lean](#), lines 13–49 first obtains height models from the all-density model theorem, then invokes the Bohr-model Stage 13.6 constructor with the same integer m and the same budgets. The dense normalization is handled before selecting R ; it does not change the prescribed partition size.

The Bohr-model constructor in [Proofs13CommonStepSelection.lean](#), lines 212–236 invokes the short-difference-rule constructor. In that file the common-step row selection, lines 92–132, explicitly returns a proper R with $\ell R = m$ or $m + 1$, as well as its step, linearity and mass facts. The short-difference-rule proof obtains this very R and the length alternative at lines 152–154, forms its Stage 13.6 record using the same R at line 160, and verifies all stage conjuncts at lines 166–181. One may therefore retain the length alternative together with the same proved conjuncts. None of these steps assumes $a \leq 1/6$.

The adaptive cap in the dense height model affects the preliminary fibres and their corrected mass; it does not alter the subsequent common-step partition into lengths m or $m + 1$. Thus (II.197.12) is available for suitably chosen outputs of the full all-density construction.

This is a proof-level strengthened existence statement. The weaker exported type, which records only $\ell R \leq \ell Q_{\text{st}}$, does not imply shortness for an arbitrary old witness. We never make that inference.

[write] **7 October 2026; sources 47–53.** *Changed since the pin* (checked at HEAD 06b30ca52). These wrappers were revised in 7feaa1604: `lemma_13_6_from_initial_stages_all_densities_with_length` ([Proofs13AllDensityExtraction.lean](#), line 13) now exports $m \leq \ell R \leq \ell Q_{\text{st}}$, and `common_step_row_selection` ([Proofs13CommonStepSelection.lean](#), line 92) still returns $\ell R \in \{m, m + 1\}$; the upper bound $\ell R \leq m + 1$ used here is still not exported, so this section’s proof-level argument is needed, and Question 1 of Section V.202 stands. The line numbers above refer to the pin.

II.201 Short parents and direct square geometry

Source 53, delivered Section 7.

This section is self-contained apart from the following explicitly stated analytic input.

Imported all-scale coefficient localization. Let N be prime, T a proper nonzero-step progression of positive length v , and X_0 a nonempty column set. Suppose $0 < \delta \leq 1$, $J \subseteq T$, A_c, B_c are scalar order-eight Freiman maps on J , and

$$C_0 \subseteq X_0 \times (J + y), \quad |C_0| \geq \delta |X_0|v, \quad \phi(x, z) = A_c(z - y) + B_c(z - y)x \quad ((x, z) \in C_0).$$

The source theorem [extract_bilinear_cell](#) provides a proper nonzero-step progression $U \subseteq T$ and the retained set

$$D_0 = \{(x, z) \in C_0 : z - y \in U\}, \quad \ell U \geq v^e, \quad |D_0| \geq \delta |X_0| \ell U, \quad e = \delta^2 / 2^{14}, \quad (\text{II.201.1})$$

on which ϕ is bilinear. There is no lower-scale premise such as $8 \leq v^e$. The exponent is the one-domain simultaneous coefficient exponent; the two scalar maps do not cause two successive exponent losses. In our application,

$$\delta = 2^{-135} a^{704}, \quad e = 2^{-284} a^{1408} = g(a). \quad (\text{II.201.2})$$

Lemma II.201.1 (Rectification inside a short progression). *Let P be a proper progression in $\mathbb{F}_N$ of nonzero step d and length n , with $2(n - 1) < N$. If a proper progression X of length at least two is contained in P , then, after possibly reversing only its indexing, there is a positive integer s such that*

$$\text{step } X = sd, \quad s(\ell X - 1) \leq n - 1. \quad (\text{II.201.3})$$

Its carrier and length are unchanged by the reversal.

Proof. Normalize by the affine bijection $z \mapsto (z - P.\text{start})/d$. The successive points of X have unique indices $i_j \in \{0, \dots, n - 1\}$. All differences $i_{j+1} - i_j$ lie in $[-(n - 1), n - 1]$ and represent the same residue step X/d . Their pairwise differences have absolute value at most $2(n - 1) < N$, so they are equal as integers. Let the common value be s_0 . It is nonzero because X is proper and has at least two points. If it is negative, reverse the indexing, and put $s = |s_0|$. Comparing extreme indices proves (II.201.3). $\square$

Lemma II.201.2 (Balanced aligned parents). *If integers $L \geq 2$, $s \geq 1$ and q_Q satisfy $s(L - 1) < q_Q$, then the index interval $\{0, \dots, q_Q - 1\}$ partitions into step- s progressions with lengths between $\lceil L/2 \rceil$ and L .*

Proof. Split the interval into its s residue chains. Each has length $n \geq L - 1$. For a chain of length n , put $c = \lceil n/L \rceil$, and divide it into c consecutive blocks whose lengths differ by at most one. They have length at most L . If $c = 1$, their length is at least $L - 1 \geq \lceil L/2 \rceil$. If $c \geq 2$, then

$$n \geq (c - 1)L + 1, \quad n/c \geq L - (L - 1)/c \geq (L + 1)/2.$$

Thus the smallest block has length at least $\lfloor (L + 1)/2 \rfloor = \lceil L/2 \rceil$. These blocks partition each chain and hence the whole interval. $\square$

The lower bound is sharp for this partition class: a chain of length $L + 1$ needs at least two blocks of length at most L , and one has length at most $\lceil L/2 \rceil$.

Lemma II.201.3 (Endpoint-contained common-step cover). *Let $S = P(x, d; L)$ and $T = P(y, e; M)$ be proper progressions in a cyclic group, with $e \neq 0$ and $e = jd$ for a positive integer j . If an integer k satisfies*

$$1 \leq k \leq M, \quad jk \leq L, \quad (\text{II.201.4})$$

and $D \subseteq S \times T$ has $|D| \geq \delta LM$ with $\delta \geq 0$, there are proper progressions $V \subseteq S$ and $W \subseteq T$ of common step e and length k such that

$$|D \cap (V \times W)| \geq (\delta/4)k^2.$$

Every restriction-invariant property of D , including agreement with a fixed bilinear function, persists on the intersection.

Proof. Split the indices of S into their j residue chains. Their lengths are at least k by $jk \leq L$, their steps are $jd = e$, and their total length is L . In any chain of $n \geq k$ points, take consecutive blocks of length k , then add the terminal length- k block if there is a nonempty remainder. All blocks stay in the chain and cover it. Their total length is

$$k \lceil n/k \rceil < n + k \leq 2n.$$

Thus S has a contained step- e cover with total capacity less than $2L$. Apply the same block construction to T to get capacity less than $2M$.

The product cells cover $S \times T$, with possible overlap, and their total capacity is less than $4LM$. Counting members of D with multiplicity yields at least $|D| \geq \delta LM$. Some cell therefore has density at least $\delta/4$. Each cell is proper by containment in an injective parent indexing, and has the required common step and side. Restriction proves the final assertion. $\square$

Proposition II.201.4 (Short-parent extraction with the full exponent). *Let N be prime and let R, Q_{st}, S be proper nonzero-step progressions with*

$$\text{step } R = \text{step } Q_{\text{st}}, \quad S \subseteq R, \quad L = \ell S \geq 3, \quad r = \ell R \leq \ell Q_{\text{st}}, \quad 2(r-1) < N.$$

Suppose the support, mass and row identities in (II.197.14) hold for some $0 < \delta \leq 1$, with scalar order-eight Freiman coefficients on $J_0 \subseteq Q_{\text{st}}$. Then there is a proper common-step square, with retained set $B \subseteq C$, density at least $\delta/4$, bilinearity of the original ϕ on B , and side

$$k \geq \lceil L/2 \rceil^{\delta^2/2^{14}} - 1. \quad (\text{II.201.5})$$

Proof. Rectify $S \subseteq R$ by Lemma II.201.1, reversing its indexing if needed. Then

$$\text{step } S = s \text{step } Q_{\text{st}}, \quad s \geq 1, \quad s(L-1) \leq r-1 < \ell Q_{\text{st}}.$$

Transport the partition of Lemma II.201.2 through the proper indexing of Q_{st} . It yields disjoint parents T_i of step $\text{step } S$ and lengths v_i with $\lceil L/2 \rceil \leq v_i \leq L$ and $\sum_i v_i = \ell Q_{\text{st}}$. The sets $C_i = \{(x, z) \in C : z - y \in T_i\}$ partition C . Weighted averaging of their actual pair counts therefore selects a parent T of length v with

$$|C_T| \geq \delta Lv, \quad C_T \subseteq S \times ((J_0 \cap T) + y).$$

The row identity and Freiman maps restrict to $J_0 \cap T$. Apply the analytic input (II.201.1) with this same density δ . It gives $U \subseteq T$, length $M \geq v^e$ for $e = \delta^2/2^{14}$, and a set $D_0 \subseteq C_T \cap (S \times (U + y))$

with $|D_0| \geq \delta LM$ and bilinearity on D_0 . No density is lost in the parent selection or in this analytic localization.

Since $L \geq 3$, $v \geq 2$ and $e > 0$, so the integer M is at least two. Further, $2(v-1) \leq 2(L-1) \leq 2(r-1) < N$. Rectify $U \subseteq T$ with a possible indexing reversal. For a positive integer j this gives

$$\text{step } U = j \text{ step } S, \quad j(M-1) \leq v-1 \leq L-1.$$

Set $k = M-1$. It is positive, $k \leq M$ and $jk \leq L$. Apply Lemma II.201.3 to S and $U+y$, retaining a set $B \subseteq D_0$ of mass at least $(\delta/4)k^2$. Finally

$$k = M-1 \geq v^e - 1 \geq \lceil L/2 \rceil^e - 1.$$

Properness, equal nonzero step, original-domain containment and bilinearity have all been preserved. $\square$

The original phase and the offset. Only carriers are restricted, or their indexing reversed. Neither operation changes the ambient variables in ϕ . More explicitly, if the localized affine coefficients are $A_c(h) = a_1h + a_0$ and $B_c(h) = b_1h + b_0$, substituting $h = z - y$ gives

$$\phi(x, z) = (a_0 - a_1y) + (b_0 - b_1y)x + a_1z + b_1xz.$$

This includes both offset corrections and is a bilinear formula in the original coordinates. There is no claim that the old arbitrary Stage 13.9 witness acquired a grid; the selected U is newly constructed inside a short aligned parent.

II.202 The sharper coefficient and final assembly

Source 53, delivered Section 8.

For $N \geq N_0(a_0)$, Proposition II.199.1 implies all seven budgets. Proposition II.198.1 and the verified short-parent construction provide every hypothesis of Proposition II.201.4. In particular, $L \geq W(a) \geq 2^{135}$, so no singleton or two-point exception is needed.

Use $\delta = 2^{-135}a^{704}$ and $g(a) = \delta^2/2^{14}$. With $A_* = c(a)N^{\beta(a)}$, (II.201.5) and $L \geq A_*^{f(a)}$ give

$$\begin{aligned} k &\geq \lceil L/2 \rceil^{g(a)} - 1 \geq (L/2)^{g(a)} - 1 \\ &\geq 2^{-g(a)}c(a)^{f(a)g(a)}N^{\beta(a)f(a)g(a)} - 1 = D(a)N^{4\gamma(a)} - 1. \end{aligned} \quad (\text{II.202.1})$$

The full $\lceil L/2 \rceil^{g(a)}$ estimate is retained. Replacing it by the weaker $L^{g(a)/2}$ bound would unnecessarily halve this exponent.

Lemma II.202.1 (Uniform coefficient bound). *For every $0 < a \leq 1$, $D(a) > 1/2$.*

Proof. Let $h = \log_2(1/a) \geq 0$. Since $\log_2(1/c(a)) = 1 + 2^{114}a^{-320}(228 + 576h)$, the exact loss is

$$\begin{aligned} \log_2(1/D(a)) &= g(a) + f(a)g(a)\log_2(1/c(a)) \\ &= 2^{-284}a^{1408} + 2^{-384}a^{1856} + 2^{-270}a^{1536}(228 + 576h). \end{aligned} \quad (\text{II.202.2})$$

Using $\log 2 \geq 1/2$ and $e^x \geq 1 + x$, $2^{1536h} \geq 1 + 768h$. Hence

$$228 + 576h \leq 228(1 + 768h) \leq 228 \cdot 2^{1536h}, \quad a^{1536}(228 + 576h) \leq 228.$$

The loss in (II.202.2) is consequently at most

$$2^{-284} + 2^{-384} + 228 \cdot 2^{-270} < 230 \cdot 2^{-270} < 256 \cdot 2^{-270} = 2^{-262} < 1.$$

Thus $\log_2(1/D(a)) < 1$, proving $D(a) > 1/2$ uniformly, including $a = 1$ and arbitrarily small positive a . $\square$

Completion of Theorem F.58.1. Equation (II.202.1) and Lemma II.202.1 give the first two side estimates. Put $x = N^{\gamma(a)}$. Monotonicity of γ and $N \geq R_7$ imply

$$x \geq N^{\gamma(a_0)} \geq 8/c(a_0) \geq 16.$$

Therefore

$$\frac{1}{2}x^4 - 1 \geq 8x^3 - 1 \geq x^3 = N^{3\gamma(a)}.$$

The cover retains density $\delta/4 = 2^{-137}a^{704}$ and a subset of the original A ; all phase identities use the original ϕ . Properness gives $|V| = |W| = k$. This proves every conclusion of the theorem. The previously established comparison $N_0 \leq N_{\text{safe}}$ now proves Corollary F.58.2 as well. $\square$

II.203 The exact retained interface

Source 51, delivered Section 2.

We first translate the source interface into mathematical notation. This also fixes all normalization and indexing conventions used in the proof.

II.203.1 The context and the two initial stages

An order-eight Freiman homomorphism on a set preserves equal sums of eight elements. A function $\phi : (\mathbb{Z}/N\mathbb{Z})^2 \rightarrow \mathbb{Z}/N\mathbb{Z}$ is separately order-eight Freiman on A when each of its vertical and horizontal restrictions has this property on the corresponding section of A .

An eight-arrangement at height h consists of sixteen horizontal coordinates x_i and sixteen vertical coordinates y_i , with

$$\sum_{i=1}^8 x_i = \sum_{i=9}^{16} x_i, \quad (x_i, y_i), (x_i, y_i + h) \in A \quad (1 \leq i \leq 16).$$

It is respected when the analogous equality holds for the sixteen values $\phi(x_i, y_i + h) - \phi(x_i, y_i)$. Let $C(h)$ count all such arrangements and let $C_\phi(h)$ count the respected ones. A context consists of A, ϕ, α, η satisfying

$$0 < \alpha \leq 1, \quad |A| = \alpha N^2, \quad \eta = 2^{-44}, \quad (1 - \eta) \sum_h C(h) \leq \sum_h C_\phi(h),$$

together with the separate Freiman property. A height is good when $(1 - 2\eta)C(h) \leq C_\phi(h)$ and strong when it is good and $C(h) \geq \alpha^{32}N^{31}/16$.

Let

$$\mathcal{E}_h = \{(x, y) : (x, y), (x, y + h) \in A\}, \quad f_h(x) = |\{y : (x, y) \in \mathcal{E}_h\}|.$$

The source Fourier transform is unnormalized: $\widehat{f}(\xi) = \sum_x f(x) \exp(\pm 2\pi i \xi x / N)$; its sign convention does not affect any norm below. Write $\|\xi\|_N$ for the absolute value of a centered integer representative. For $\theta > 0$ put

$$\theta_1 = 2^{-1882} \theta^{10477}, \quad Q_{134}(\theta) = 2^{1882} \theta^{-10479}.$$

Stage 13.4 stores a positive integer q , an integer m_D , a progression P , a set H_{134} , and q pairs $a_i, b_i \in \mathbb{Z}/N\mathbb{Z}$. Its complete conditions are:

134.1 $d_P \neq 0$, P proper, $H_{134} \subseteq P$, and $q \leq Q_{134}(\theta)$;

134.2 $m_D = \lfloor \theta_1 N^{\theta_1^2/(16q)} / (64\pi) \rfloor$, and either $\ell(P) = m_D$ or $\ell(P) + 1 = m_D$;

134.3 $\alpha^{32} N^{31} \ell(P) / 8 \leq \sum_{h \in H_{134} \cap \mathcal{H}_{\text{good}}} C(h)$;

134.4 if $h \in H_{134}$ and $|\widehat{f}_h(\xi)| \geq \theta N^2$, then $\xi = a_i h + b_i$ for some i .

Here the floor is the exact predicate $m_D \leq x < m_D + 1$. Stage 13.5 stores Q and defines

$$\mathcal{H} = (Q \cap H_{134}) \cap \{h : h \text{ is strong}\}.$$

Its complete conditions are $d_Q \neq 0$, properness, $Q \subseteq P$, and

$$\frac{1}{2}(\ell(P))^{2^{-12q}} \leq \ell(Q), \tag{II.203.1}$$

$$\|(a_i h + b_i) d_Q\|_N \leq (\ell(P))^{-2^{-11q}} N \quad (i < q, h \in Q), \tag{II.203.2}$$

$$\alpha^{32} \ell(Q) / 20 \leq |\mathcal{H}|. \tag{II.203.3}$$

II.203.2 The six numerical budgets

The Section 13 scalar functions and the relevant Section 10 functions are

$$\begin{aligned} K_{13}(\alpha) &= 2^{114} \alpha^{-320}, & \zeta_{13}(\alpha) &= 2^{-228K_{13}(\alpha)} \alpha^{576K_{13}(\alpha)}, \\ Q_{13}(\alpha) &= 2^{220} \alpha^{-221}, & \lambda_{10}(a) &= 2^{-37} a^{11/2}, \\ k(a) &= \lceil 2^{74} a^{-10} \rceil, & \zeta_{10}(a) &= 2^{-155k(a)} a^{18k(a)} / k(a). \end{aligned}$$

In particular, the ceiling and the division by $k(a)$ in the last formula must both be retained. Define

$$\begin{aligned} c &= \zeta_{13}(\alpha) / 2, & \beta &= 2^{-13Q_{13}(\alpha)}, \\ f &= 2^{-100} \alpha^{448}, & W &= 2^{135} \alpha^{-704}, & z &= \zeta_{10}(\alpha^{32}/16). \end{aligned} \tag{II.203.4}$$

Besides N prime, the context, $\alpha \leq 1/6$, and the two complete initial-stage predicates at $\theta = \lambda_{10}(\alpha^{32}/16)$, the budget theorem assumes an integer m satisfying exactly

$$0 < m, \quad m^2 \leq N, \quad m + 1 \leq \ell(Q), \tag{II.203.5}$$

$$cN^\beta \leq m, \quad (\ell(P))^{-2^{-11q}} \leq z/m, \quad W \leq (cN^\beta)^f. \tag{II.203.6}$$

These are, in order, **hm**, **hsize**, **hupper**, **hlower**, **hbudget**, and **hwidth**. The integer m in this theorem is *not* the field m_D stored in Stage 13.4. In our construction they will be n and p , respectively.

The theorem then existentially produces F, G, H, J , with the four subsequent stage predicates and $\ell(R) \leq \ell(Q)$. Section II.208 lists every conjunct of these predicates while verifying our explicit choices. This avoids treating a weakened or obsolete version of Stage 13.7 as the current interface.

II.203.3 The grid is a stronger assertion

The source predicate `Corollary1310Grid` asks for a positive finite number of product cells $V_i \times W_i$ which form a disjoint partition of all $S \times (U + y)$. For every i it requires

$$d_{V_i} \neq 0, \quad d_{V_i} = d_{W_i}, \quad V_i, W_i \text{ proper}, \quad \ell(V_i) = \ell(W_i) \geq \sqrt{\ell(U)} - 1.$$

Different cells may have different common steps. The counterexample will exclude even one cell of side length two, whatever that common step might be. It therefore does not depend on the extra difficulty of demanding an exact disjoint partition.

II.204 The geometric obstruction in isolation

Source 51, delivered Section 3.

Lemma II.204.1 (Separated difference sets). *Let $n \geq 2$ and $N > n^2 - 1$. In $\mathbb{Z}/N\mathbb{Z}$ put*

$$S_n = \{0, 1, \dots, n-1\}, \quad U_n = \{0, n, 2n, \dots, (n-1)n\}.$$

Then $(S_n - S_n) \cap (U_n - U_n) = \{0\}$. Hence no proper common-step square of side at least two is contained in $S_n \times U_n$.

Proof. A member of the intersection has representatives a and nj with $|a| \leq n-1$ and $|j| \leq n-1$. If $a \equiv nj \pmod{N}$, then

$$|a - nj| \leq (n-1) + n(n-1) = n^2 - 1 < N.$$

Thus $a = nj$ as integers. Since $|a| < n$, this forces $a = j = 0$. If $V \subseteq S_n$ and $W \subseteq U_n$ are progressions with the same step e and at least two terms, their first two terms put e in both difference sets. The common step must therefore be zero, contrary to the required nonzero step. $\square$

For a small illustration take $N = 47$ and $n = 5$:

$$S = \{0, 1, 2, 3, 4\}, \quad U = \{0, 5, 10, 15, 20\}.$$

The two nonzero difference sets are $\{\pm 1, \dots, \pm 4\}$ and $\{\pm 5, \pm 10, \pm 15, \pm 20\}$; they are disjoint modulo 47. With $R = S$ and $Q = \{0, \dots, 20\}$, one has $S \subseteq R$, $U \subseteq Q$, $|R| \leq |Q|$, and $d_U = 5d_S$. Since $\sqrt{5} - 1 > 1$, this also excludes the grid predicate. This example is an illustration of the geometry only: it is not presented as a full-budget witness. The main task is to embed the same mechanism into genuine initial-stage data at the much larger scales demanded by (II.203.5)–(II.203.6).

II.205 Uniform constants and a simultaneous prime choice

Source 51, delivered Section 4.

All constants introduced before choosing N in this section are fixed independently of N and its eventual density. They are real numbers or integers defined symbolically; numerical expansion is neither needed nor intended.

II.205.1 The fixed exponents

Set

$$\begin{aligned} K &= 2^{762}, & q &= 2K + 1 = 2^{763} + 1, & T &= 1882 + 763 \cdot 10477 = 7\,995\,833, \\ d_0 &= \frac{2^{-T}}{64\pi}, & c_P &= \frac{d_0}{2}, & u_0 &= \frac{2^{-2T}}{16q}, & w &= 2^{-11q}, & r &= \frac{u_0 w}{4}. \end{aligned} \quad (\text{II.205.1})$$

We use T for the fixed exponent to keep it distinct from an integer step multiplier later. Uniform bounds for the Section 10 and 13 functions will use

$$\begin{aligned} k_0 &= 2^{1394}, & c_0 &= 2^{-(2532k_0+1)}, & z_0 &= 2^{-(2531k_0+1394)}, \\ \beta_0 &= 2^{-13 \cdot 2^{9437184}}, & \beta_1 &= 2^{-13 \cdot 2^{7340032}}, & f_0 &= 2^{-1892}, & W_0 &= 2^{2951}. \end{aligned} \quad (\text{II.205.2})$$

These extremely small or large quantities are finite and positive. They will occur only in inequalities and fixed thresholds.

Lemma II.205.1 (The exponent margin). *One has $0 < \beta_1 < r < u_0/4 < 1/64$, $0 < w \leq 1/2$, and $u_0 - 2r > 0$.*

Proof. Since $q < 2^{764}$,

$$r = \frac{2^{-2T-11q}}{64q} > 2^{-(2T+770+11q)}.$$

The exact integer identity $2T + 770 = 15\,992\,436 < 2^{25}$ and the bound $11q < 2^{768}$ imply $2T + 770 + 11q < 2^{769}$. Therefore

$$r > 2^{-2^{769}} > 2^{-13 \cdot 2^{7340032}} = \beta_1.$$

The remaining inequalities follow from $q \geq 1$, $0 < 2^{-2T} < 1$, $0 < w < 1$, and $u_0 - 2r = u_0(1 - w/2) > 0$. $\square$

II.205.2 Uniformity over the actual density interval

Lemma II.205.2 (Uniform scalar estimates). *For every $\alpha \in [1/16, 1/8]$, the notation in (II.203.4) satisfies*

$$c_0 \leq c \leq \frac{1}{2}, \quad \beta_0 \leq \beta \leq \beta_1 < r, \quad f_0 \leq f < 1, \quad W \leq W_0, \quad z \geq z_0 > 0. \quad (\text{II.205.3})$$

Also, with $\theta = \lambda_{10}(\alpha^{32}/16)$ and $\theta_1 = 2^{-1882}\theta^{10477}$,

$$\theta = 2^{-59}\alpha^{176}, \quad 2^{-763} \leq \theta \leq 2^{-587} < 1, \quad 0 < \theta \leq \alpha^{32}/4, \quad \theta_1 \geq 2^{-T}. \quad (\text{II.205.4})$$

Proof. The identity for θ is $2^{-37}(\alpha^{32}/16)^{11/2} = 2^{-59}\alpha^{176}$. The endpoints of the density interval give the powers -763 and -587 . Moreover $4\theta/\alpha^{32} = 2^{-57}\alpha^{144} \leq 1$, and the definition of T yields the bound on θ_1 .

For the Section 13 prefactor, $K_{13}(\alpha) \leq 2^{114+4 \cdot 320} = k_0$. Since $\alpha \geq 2^{-4}$,

$$\zeta_{13}(\alpha) = 2^{-228K_{13}(\alpha)}\alpha^{576K_{13}(\alpha)} \geq 2^{-2532K_{13}(\alpha)} \geq 2^{-2532k_0}.$$

This gives $c \geq c_0$. Positivity and $\alpha \leq 1$ give $0 < c \leq 1/2$. Next,

$$2^{7340032} \leq Q_{13}(\alpha) \leq 2^{9437184},$$

because $2^{20} + 3 \cdot 2^{21} = 7\,340\,032$ and $2^{20} + 4 \cdot 2^{21} = 9\,437\,184$. Taking the decreasing function $x \mapsto 2^{-13x}$ gives the asserted bounds for β . The formulas for f_0, W_0 follow from $100 + 4 \cdot 448 = 1892$ and $135 + 4 \cdot 704 = 2951$.

Finally put $a = \alpha^{32}/16$. Then $2^{-132} \leq a \leq 2^{-100}$ and $2^{74}a^{-10} \leq 2^{1394} = k_0$. Because k_0 is an integer, the ceiling $k = \lceil 2^{74}a^{-10} \rceil$ satisfies $1 \leq k \leq k_0$, without an extra $+1$ loss. Retaining the division by k ,

$$z = 2^{-155k} a^{18k} / k \geq 2^{-(155+18 \cdot 132)k} / k \geq 2^{-2531k_0} / k_0 = z_0.$$

This argument uses no continuity across jumps of the ceiling. $\square$

II.205.3 One threshold for all subsequent requirements

Choose a prime N large enough that the following eight conditions hold:

$$N \geq K^2, \quad d_0 N^{u_0} \geq 2, \quad (\text{II.205.5})$$

$$c_P N^{u_0-2r} \geq 1, \quad N^r \geq 5, \quad (\text{II.205.6})$$

$$N^{1-r} \geq 16, \quad N^{1-2r} > 2, \quad (\text{II.205.7})$$

$$c_P^{-w} N^{-3r} \leq z_0, \quad c_0 N^{\beta_0} \geq W_0^{1/f_0}. \quad (\text{II.205.8})$$

Every constant here is fixed. All growth exponents are positive by Lemma II.205.1; the first inequality in (II.205.8) is equivalently $N^{3r} \geq c_P^{-w}/z_0$. Thus all eight conditions hold for every real N above one finite threshold. There are arbitrarily large primes beyond it.

Now set

$$\begin{aligned} b &= \lfloor N/8 \rfloor, & \alpha &= b/N, & \theta &= 2^{-59} \alpha^{176}, & \theta_1 &= 2^{-1882} \theta^{10477}, \\ d &= \theta_1/(64\pi), & u &= \theta_1^2/(16q), & p &= \lfloor dN^u \rfloor, & n &= \lfloor N^r \rfloor. \end{aligned} \quad (\text{II.205.9})$$

The earlier density bounds apply because $N \geq K^2 \geq 16$: $b \geq N/8 - 1 \geq N/16$, while $b \leq N/8$. Hence $1/16 \leq \alpha \leq 1/8 < 1/6$. In particular $d \geq d_0$, $u \geq u_0$, $d < 1$, and $0 < u < 1/(16q) \leq 1/16$.

Lemma II.205.3 (Length and containment estimates). *The chosen integers satisfy*

$$\begin{aligned} p &\geq c_P N^{u_0} \geq N^{2r} \geq n^2, & 1 &\leq p < N, \\ 5 &\leq n \leq N^r \leq N/16 \leq b, & n &\geq N^r/2, & 2n^2 &< N. \end{aligned} \quad (\text{II.205.10})$$

Proof. The inequality $dN^u \geq d_0 N^{u_0} \geq 2$ permits $\lfloor dN^u \rfloor \geq dN^u/2 \geq c_P N^{u_0}$. The third threshold condition then gives $p \geq N^{2r}$. Since $d < 1$, $u < 1$, and $N > 1$, we have $p \leq dN^u < N$. The remaining assertions follow from $N^r \geq 5$, the elementary floor bound $\lfloor x \rfloor \geq x/2$ for $x \geq 2$, and (II.205.7). $\square$

This is the entire large-modulus argument. In particular, it does not invoke a theorem with a fixed-density threshold and then substitute a moving density into it. The uniform estimates prove a single threshold valid for our actual $\alpha(N)$.

II.206 A genuine strip context and the initial data

Source 51, delivered Section 5.

Let $X = \mathbb{Z}/N\mathbb{Z}$, let $B_0 = \{0, \dots, b-1\} \subset X$, and put

$$A = B_0 \times X, \quad \phi(x, y) = 0, \quad \eta = 2^{-44}.$$

The b residues in B_0 are distinct, so $|A| = bN = \alpha N^2$. The zero map preserves every additive relation on every horizontal or vertical section. It respects every arrangement. These facts verify each field of `Section13Context`, including the global respected-arrangement proportion.

II.206.1 The fixed-height arrangement count

Let $\rho(s)$ be the number of ordered eight-tuples of elements of B_0 having sum $s \in X$, and write

$$E_8(B_0) = \sum_{s \in X} \rho(s)^2.$$

Since $\sum_s \rho(s) = b^8$, Cauchy–Schwarz gives

$$E_8(B_0) \geq b^{16}/N = \alpha^{16} N^{15}. \quad (\text{II.206.1})$$

For a prescribed height h , an arrangement has an additive sixteen-tuple of horizontal coordinates in B_0 , and its sixteen vertical coordinates can be chosen freely and independently. Therefore the exact normalization is

$$C(h) = C_\phi(h) = N^{16} E_8(B_0) \geq \alpha^{16} N^{31}. \quad (\text{II.206.2})$$

There is no extra factor N at fixed height. Summing over heights would introduce one. Every height is good, and every height is strong because $\alpha^{16} \geq \alpha^{32}/16$.

II.206.2 A uniformly covering family of constant frequencies

At every height, $\mathcal{E}_h = A$ and $f_h = N\mathbf{1}_{B_0}$. For a nonzero frequency ξ , let $j = \|\xi\|_N$, so $1 \leq j \leq N/2$. The finite geometric sum gives

$$\left| \sum_{x=0}^{b-1} e^{\pm 2\pi i \xi x/N} \right| = \frac{|\sin(\pi b \xi/N)|}{|\sin(\pi \xi/N)|} \leq \frac{1}{\sin(\pi j/N)} \leq \frac{N}{2j}.$$

The last inequality is $\sin t \geq 2t/\pi$ on $[0, \pi/2]$. Thus

$$|\widehat{f}_h(\xi)| \leq N^2/(2j). \quad (\text{II.206.3})$$

It follows that $|\widehat{f}_h(\xi)| \geq \theta N^2$ forces $j \leq 1/(2\theta) \leq 2^{762} = K$. Zero frequency is included separately; no division by zero is used.

Index the integer interval $[-K, K]$ by $i = 0, \dots, q-1$, and set $a_i = 0$, $b_i = i - K \pmod{N}$. All Fourier coefficients at or above the threshold are on these constant graphs, for every height. The number q is fixed independently of N , and

$$q < 2^{764} < 2^{1882} \leq 2^{1882} \theta^{-10479} = Q_{134}(\theta). \quad (\text{II.206.4})$$

All listed frequencies have centered absolute value at most K , including frequencies that are unnecessary for covering the actual large spectrum.

II.206.3 Every Stage 13.4 and 13.5 condition

Take $P = Q = P(0, 1; p)$ and $H_{134} = P$, with $m_D = p$ and the graph family above. Properness and nonzero step follow from $p < N$ and $N > 1$. The exact floor equation is the definition of p in (II.205.9); the first branch $\ell(P) = m_D$ of the source's length alternative holds. Since every height is good, (II.206.2) yields

$$\sum_{h \in H_{134} \cap \mathcal{H}_{\text{good}}} C(h) = pN^{16} E_8(B_0) \geq p\alpha^{16} N^{31} \geq \alpha^{32} N^{31} p/8.$$

Together with $q > 0$, (II.206.4), and the Fourier cover, this checks all four Stage 13.4 groups listed in Section II.203. The cutoff also satisfies $0 < \theta \leq \alpha^{32}/4$, the extra cutoff hypothesis in the separate statement of Lemma 13.4.

For Stage 13.5, containment $Q \subseteq P$ is equality. Since $p \geq 1$ and $0 < 2^{-12q} \leq 1$, inequality (II.203.1) follows from $p^{2^{-12q}}/2 \leq p$. For every listed graph and every $h \in Q$,

$$\|(a_i h + b_i) d_Q\|_N \leq K \leq \sqrt{N} \leq N^{1-w} \leq p^{-w} N.$$

Here $N \geq K^2$, $w \leq 1/2$, and $p \leq N$ were used in that order. Thus the frequency condition is verified for the whole graph family, rather than just for the graphs needed at a particular height. Finally, every height is strong and $H_{134} = Q$, so

$$\mathcal{H} = Q, \quad |\mathcal{H}| = p, \quad \alpha^{32} p/20 \leq p. \quad (\text{II.206.5})$$

This completes Stage 13.5 without omitting any frequency or mass condition.

II.207 Verification of all six budgets

Source 51, delivered Section 6.

Set the theorem's independent integer input to $m = n$. The first three budgets follow at once from Lemma II.205.3:

$$n \geq 5 > 0, \quad n^2 < N/2 < N, \quad n + 1 \leq n^2 \leq p = \ell(Q).$$

For the lower-length budget, $c \leq 1/2$, $\beta < r$, $N \geq 1$, and $n \geq N^r/2$ give

$$cN^\beta \leq \frac{1}{2} N^r \leq n. \quad (\text{II.207.1})$$

Thus the floor is absorbed uniformly; no extra density-dependent threshold is hidden in this step.

The corrected Bohr budget uses the whole progression P and the radius z , including its denominator from Section 10. Since $u_0 w = 4r$,

$$\begin{aligned} np^{-w} &\leq N^r (c_P N^{u_0})^{-w} \\ &= c_P^{-w} N^{r-u_0 w} = c_P^{-w} N^{-3r} \leq z_0 \leq z. \end{aligned}$$

Dividing by $n > 0$ is exactly $p^{-w} \leq z/n$.

For the width budget it is essential to use the lower target cN^β , not just the actual chosen integer n . Uniform positivity and (II.205.8) give

$$cN^\beta \geq c_0 N^{\beta_0} \geq W_0^{1/f_0} \geq 1.$$

Because $f \geq f_0 > 0$, raising this base to f gives

$$(cN^\beta)^f \geq (cN^\beta)^{f_0} \geq W_0 \geq W. \quad (\text{II.207.2})$$

This verifies the source's actual **hwidth**, a stronger check than merely arranging a wide enough final progression. Together with $\alpha < 1/6$ and Section II.206, every antecedent of the budget-qualified extraction theorem is now satisfied.

II.208 The complete chosen output chain

Source 51, delivered Section 7.

We next construct the retained output data directly. Set

$$\begin{aligned} R = S = P(0, 1; n), \quad U = P(0, n; n), \quad y = 0, \\ Y(h) = A \quad (h \in X), \quad B = S \times Q, \\ a(h) = c(h) = 0 \quad (h \in X), \quad J_{138} = Q, \quad C = S \times Q, \quad D_{139} = S \times U. \end{aligned} \quad (\text{II.208.1})$$

The letters a, c on the middle line denote row coefficients, not the scalar c in (II.203.4). We use J_{138}, D_{139} for sets so that the Stage 13.8 and 13.9 record names H, J do not collide with set notation. The source fields are respectively $H.J$ and $J.D$.

The progressions R, S are proper and have nonzero step one. The integer representatives of U are $0, n, \dots, (n-1)n$, all below $n^2 < N$; hence U is proper and has nonzero step n . Also

$$S = R \subseteq B_0, \quad U \subseteq Q, \quad \ell(R) = n \leq p = \ell(Q). \quad (\text{II.208.2})$$

In particular, $U \subseteq Q$ follows from $(n-1)n < n^2 \leq p$. None of these containment statements relies on a modular quotient estimate.

II.208.1 Stage 13.6

The source requires: the step equality $d_R = d_Q$, nonzero step, properness, the length bound $cN^\beta \leq \ell(R)$, and for every $h \in \mathcal{H}$,

$$Y(h) \subseteq \mathcal{E}_h, \quad |Y(h)| \geq 2^{-43} \alpha^{224} N^2,$$

with the vertical ϕ -difference affine in the first coordinate on $Y_R(h) = \{(x, y) \in Y(h) : x \in R\}$. It additionally requires the two separate mass inequalities

$$2^{-43} \alpha^{224} \ell(R) N |\mathcal{H}| \leq \sum_{h \in \mathcal{H}} |Y_R(h)|, \quad (\text{II.208.3})$$

$$2^{-48} \alpha^{256} \ell(Q) \ell(R) N \leq \sum_{h \in \mathcal{H}} |Y_R(h)|. \quad (\text{II.208.4})$$

The step and length conditions follow from (II.207.1) and the definitions. Here $Y(h) = A = \mathcal{E}_h$, and $|Y(h)| = \alpha N^2 \geq 2^{-43} \alpha^{224} N^2$ because $2^{-43} \alpha^{223} \leq 1$. Since $S \subseteq B_0$, $Y_R(h) = S \times X$, of cardinality nN . The vertical difference is zero, so its affine representation has both coefficients zero. Finally (II.206.5) makes the sum exactly pnN ; both (II.208.3) and (II.208.4) follow because their scalar coefficients are at most one. The additional retained comparison $\ell(R) \leq \ell(Q)$ was verified in (II.208.2).

II.208.2 Stage 13.7 including the original domain

The complete current predicate requires $d_S \neq 0$, properness, $S \subseteq R$, and $(\ell(R))^f \leq \ell(S)$, followed by

$$\begin{aligned} B \subseteq S \times (\mathcal{H} + y), \quad 2^{-43} \alpha^{224} \ell(S) |\mathcal{H}| &\leq |B|, \\ 2^{-48} \alpha^{256} \ell(Q) \ell(S) &\leq |B|. \end{aligned}$$

Every row of ϕ on B above a critical height must be affine, and one must also have $B \subseteq A$.

All progression properties hold because $S = R$. Since $n \geq 1$ and $0 < f < 1$, $n^f \leq n$. Our support is exactly $S \times \mathcal{H}$ and $|B| = np$, so both lower mass bounds hold. Every row is zero. Crucially, $S \subseteq B_0$ gives $B = S \times Q \subseteq B_0 \times X = A$. Thus this is the full `IsStage137Data`, not the earlier `IsStage137DataWithoutDomain`.

II.208.3 Stage 13.8 with all critical heights retained

The row identity in the source is $\phi(x, y + h) = a(h) + c(h)x$ on B for $h \in \mathcal{H}$. Our zero coefficients satisfy it. The remaining conditions are

$$\begin{aligned} J_{138} \subseteq \mathcal{H}, \quad h \mapsto (a(h), c(h)) &\text{ is order-eight Freiman on } J_{138}, \\ C = \{z \in B : z_2 - y \in J_{138}\}, \quad 2^{-135} \alpha^{704} \ell(Q) \ell(S) &\leq |C|. \end{aligned}$$

They hold with $J_{138} = Q = \mathcal{H}$ and $C = B$: the coefficient pair is constant zero and $|C| = np$. Write

$$\delta = 2^{-135} \alpha^{704} \in (0, 1). \tag{II.208.5}$$

The final inequality is just $\delta np \leq np$.

Retaining all of Q in this stage is important. Replacing J_{138} prematurely by U would demand $n^2 \geq \delta np$, which is not a consequence of our scales. The later restriction to U is made only in Stage 13.9, where its density is measured relative to $S \times U$. The separate width hypothesis used to invoke the Stage 13.8 lemma also follows: $W \leq (cN^\beta)^f \leq n^f \leq n = \ell(S)$.

II.208.4 Stage 13.9

The source requires nonzero step, properness, and $U \subseteq Q$, all already verified. Its step-multiple condition

$$\exists t \in \mathbb{Z}_{\geq 0}, \quad d_U = td_S$$

is witnessed by the positive ordinary integer $t = n$. Let $g = 2^{-284} \alpha^{1408} \in (0, 1)$. The lower length condition is $(\ell(S))^g \leq \ell(U)$, which holds because $n^g \leq n$. The exact set filter in the source gives

$$\{z \in C : z_2 - y \in U \cap J_{138}\} = S \times U = D_{139}.$$

Thus $|D_{139}| = n^2 \geq \delta \ell(S) \ell(U)$, as required. Finally, a bilinear function means a polynomial $c_{00} + c_{10}x + c_{01}y + c_{11}xy$ agreeing with ϕ on the retained set. All four coefficients zero give the required witness on D_{139} . This checks every Stage 13.9 conjunct, including the exact filter identity and its density normalization.

II.209 Failure of the grid and of the new fit hypothesis

Source 51, delivered Section 8.

Lemma II.204.1 applies since $N > 2n^2$. It gives (F.56.1) and excludes every proper common-step square of side at least two inside $S \times U$.

Suppose the grid predicate held. Because $n \geq 5$, its real lower bound $\sqrt{n} - 1$ is strictly greater than one; each integral cell side length is at least two. The partition equality makes each product cell a subset of $S \times U$. Both factors of a cell are nonempty, so its first factor is a subset of S and its second factor a subset of U . Lemma II.204.1 rules this out. The grid predicate is therefore false, proving Theorem F.56.1.

The newer source theorem uses a positive integer t with

$$d_U = td_S, \quad t\lfloor\sqrt{\ell(U)}\rfloor \leq \ell(S). \quad (\text{II.209.1})$$

For our displayed progressions, every such integer step representative satisfies $t \equiv n \pmod{N}$. Since $0 < n < N$, the least positive representative is n , so every positive representative has $t \geq n$. As $\lfloor\sqrt{n}\rfloor \geq 2$,

$$t\lfloor\sqrt{n}\rfloor \geq 2n > n = \ell(S).$$

Hence the new fit condition fails for every positive representative. The slightly weaker sufficient condition with $\lfloor\sqrt{n-1}\rfloor$ also fails. This is fully consistent with the conditional cover theorem: its missing hypothesis does not follow for these witnesses.

Over a prime field, the mere existence of a natural-number multiplier is almost vacuous for nonzero steps. Indeed, for any $d, e \neq 0$, the residue e/d has a representative in $\{1, \dots, N-1\}$ and gives $e = td$. What is not automatic is the small representative bound $t \leq \ell(S)/k$ for the desired square side k . Separate length bounds, density bounds, and Bohr control at an earlier stage cannot be read as that integer bound without an additional argument.

II.210 A conditional covering repair and its exact scope

Source 51, delivered Section 9.

A finite cover suffices for dense-square extraction; an exact disjoint grid is unnecessary. This elementary observation is already reflected in the newer source. We give an endpoint-contained proof because it makes the integer fit particularly transparent, then compare it with the actual padded-cover implementation.

Proposition II.210.1 (Endpoint-contained common-step cover). *Let $S = P(a, d; L)$ and $T = P(b, e; M)$ be proper progressions in a cyclic group. Suppose $e \neq 0$, $e = td$ for a positive integer t , and let k be an integer satisfying*

$$1 \leq k \leq M, \quad tk \leq L. \quad (\text{II.210.1})$$

If $D \subseteq S \times T$ and $|D| \geq \delta LM$ with $\delta \geq 0$, there are proper progressions $V \subseteq S$, $W \subseteq T$, both of step e and length k , such that

$$|D \cap (V \times W)| \geq \frac{\delta}{4} k^2. \quad (\text{II.210.2})$$

Any property of D inherited by restriction, including agreement with a fixed bilinear function, persists on this intersection. Primality is not required.

Proof. Split the indices $0, \dots, L-1$ of S by residue modulo t . Since $tk \leq L$, all t residue chains are nonempty, and their lengths n_a are either $\lfloor L/t \rfloor$ or $\lceil L/t \rceil$, hence at least k . Each is a proper progression of step $td = e$, inherited from the injective indexing of S , and $\sum_a n_a = L$.

For a chain of $n \geq k$ points, write $n = Ak + r$ with $A \geq 1$ and $0 \leq r < k$. Take the consecutive length- k blocks starting at $0, k, \dots, (A-1)k$. If $r > 0$, append the terminal length- k block starting at $n - k$. These blocks cover the chain, stay inside it, and have total length

$$k \lceil n/k \rceil < n + k \leq 2n.$$

Apply this to every residue chain of S . The resulting step- e cover has total length $C_S < 2L$. Apply the same construction directly to T , obtaining a step- e cover with total length $C_T < 2M$.

All product cells $V_i \times W_j$ are proper common-step $k \times k$ squares, and they cover $S \times T$, possibly with overlap. Counting members of D with multiplicity gives

$$\sum_{i,j} |D \cap (V_i \times W_j)| \geq |D| \geq \delta LM, \quad \sum_{i,j} |V_i \times W_j| = C_S C_T < 4LM.$$

The weighted average density is at least $\delta LM / (C_S C_T) \geq \delta/4$, so some cell has the required density. Restriction gives the final assertion. $\square$

II.210.1 Square root rounding and the current constants

For $M \geq 2$, put $k = \lfloor \sqrt{M-1} \rfloor$. Then $k \geq 1$ and $k^2 \leq M-1 < (k+1)^2$. Since M is an integer, this implies $M \leq (k+1)^2$, and therefore

$$k \geq \sqrt{M} - 1.$$

Consequently the additional integer budget

$$\exists t \geq 1 : \quad d_U = td_S, \quad t \lfloor \sqrt{\ell(U) - 1} \rfloor \leq \ell(S) \quad (\text{II.210.3})$$

is sufficient, with $T = U + y$, for the intended side-length lower bound. Using the Stage 13.9 density $\delta = 2^{-135} \alpha^{704}$ gives exactly the weakened Corollary 13.10 coefficient $\delta/4 = 2^{-137} \alpha^{704}$. For $M = 1$, the required lower side length is zero; choosing a singleton in each nonempty ambient factor and averaging handles this case directly. For $M = 0$ the positive-length covering statement above is not applicable, and no such case is used in our counterexample.

The stronger fit $t(M-1) \leq L$ permits $k = M-1$ when $M \geq 2$. This is a conditional increase in the extracted side length, not a claim that the existing stage data imply such a fit.

II.210.2 What is already present in the newer source

The file `Proofs13CommonStepCover.lean` at the later pin contains the declarations `bilinear_square_of_integer_step_fit` and `corollary_13_10_of_integer_step_fit`. The former assumes a prime modulus, proper input progressions, $e \neq 0$, $e = td$, and precisely $t > 0$, $k > 0$, $tk \leq L$, $k \leq M$, together with support, density, and bilinearity. It constructs equal-length common-step progressions and a retained subset $E \subseteq D$ of density at least $\delta/4$ on their product.

Its cover is padded: for each residue class it uses blocks indexed through $\lfloor L/(tk) \rfloor$, with the final block allowed to extend beyond the original endpoint. The total indexed capacity is

$$t(\lfloor L/(tk) \rfloor + 1)k \leq L + tk \leq 2L.$$

The analogous capacity for the second coordinate is at most $2M$. Properness of the padded cells follows from nonzero step and $k \leq M \leq N$ in a prime modulus. Thus the output progressions need not themselves be subsets of the input progressions, but the selected set remains a subset of D . This differs from Proposition II.210.1, whose cells stay inside both original factors and whose proof works without primality.

The source’s Corollary 13.10 wrapper uses $k = \lfloor \sqrt{M} \rfloor$ and assumes (II.209.1), together with positive M , the Stage 13.9 predicate, proper S , and the support condition. It achieves the same required bound $k \geq \sqrt{M} - 1$ and quarter-density coefficient. We credit this as an existing source theorem, rather than presenting the covering repair as a missing formal result. We inspected the declaration and proof text but did not recompile it or perform a Lean axiom audit in this work.

Neither cover proves Corollary1310Grid: overlapping or padded covers are sufficient for a dense-cell argument but do not supply an exact disjoint partition. The positive result in Section II.211 obtains suitable geometry for selected witnesses by preserving a short parent and choosing a new localized height progression.

II.210.3 An existing aligned short-parent construction

A further source check at commit 2548e7fffc64de1a05eadae81b9247d9b9042c3b found `bilinear_square_of_parent_size_budget` in `Proofs13ParentLocalization.lean`. This is already a conditional coordinated-witness construction. Given coefficient data on $S \times (Q + y)$, it assumes a positive integer a with

$$d_S = ad_Q, \quad a(\ell(S))^2 \leq \ell(Q),$$

as well as the stated support, density and Freiman coefficient hypotheses and a large-scale condition $8 \leq (\ell(S) - 1)^\varepsilon$, where $\varepsilon = \text{cor711Exponent}(\delta, 1)$ is the source’s partition exponent. It constructs a same-step bilinear square of density at least $\delta/4$ and side at least $(\ell(S) - 1)^{\varepsilon/2} - 1$, with retained set contained in the original coefficient set C .

The proof partitions the height progression into short parents of length $\ell(S) - 1$ or $\ell(S)$, all with the column step, retains density by averaging, and invokes the short-parent extraction theorem. It selects new witnesses rather than proving a grid for every already retained U . The stage, grid, and common-step-cover blobs remain unchanged at this pin. Our obstruction is therefore compatible with this additional result as well. In fact our S, Q satisfy its displayed parent-size inequality with $a = 1$ because $n^2 \leq p$; this observation alone is not asserted to verify every other hypothesis of that theorem. No proof or build of this additional source result was rerun here.

II.210.4 The newer linear-span source theorem

At the later status pin 38e74b4f0a3ef56fc2402bdf5339b9d398e68d5, observed at 00:39 UTC on 7 October 2026, `Proofs13SpanLocalization.lean` already replaces the quadratic parent-size hypothesis by

$$a(\ell(S) - 1) < \ell(Q).$$

It constructs parents with lengths between $\lfloor \ell(S)/2 \rfloor$ and $\ell(S)$. Its square theorem keeps density $\delta/4$, assumes $8 \leq \lfloor \ell(S)/2 \rfloor^\varepsilon$, and gives side at least $\lfloor \ell(S)/2 \rfloor^{\varepsilon/2} - 1$. Thus linear-span localization itself is already present upstream. The additional argument in Section II.211 uses the sharper balanced lower bound $\lceil L/2 \rceil$, rectifies the output of the existing all-scale coefficient API without this large-scale premise, takes final side $\ell(U) - 1$, and recovers the short R from the original budget

construction. The next two subsections distinguish the subsequent fixed-density and uniform-density source theorems.

II.210.5 The complete fixed-density source theorem

At commit 77f7d8c1f7506254c0b2a49ce9f80ffca8d3675c, retrieved on 7 October 2026, `Proofs13CompleteSquareExtraction.lean` already contains a full context-to-square theorem. Its declaration, lines 39–48, has the precise quantifiers

$$\forall \alpha \in (0, 1] \exists N_0(\alpha) \forall N \geq N_0(\alpha) \text{ prime } \forall \text{ contexts of exact density } \alpha.$$

For each such context it constructs proper equal-length progressions V, W with the same nonzero step and a set $B \subseteq A \cap (V \times W)$ on which the original ϕ is bilinear, with

$$\begin{aligned} \ell(V) = \ell(W) &\geq N^{\gamma(\alpha)}, & |B| &\geq 2^{-137} \alpha^{704} \ell(V) \ell(W), \\ \gamma(\alpha) &= \frac{2^{-386} \alpha^{1856}}{2^{13Q_{13}(\alpha)}} = \frac{\beta f g}{4} > 0, & g &= 2^{-284} \alpha^{1408}. \end{aligned} \tag{II.210.4}$$

Its lines 13–29 define the exponent and establish this closed form. All intermediate data and square geometry are constructed in that source theorem. Therefore the actual-chain square construction is not described here as missing from the latest inspected upstream source. The stage, grid and explicit budget files were re-fetched at this pin and are byte-identical to the primary snapshots used in Theorem F.56.1.

There is a useful exact comparison with the direct finite-budget repair in Section II.211. On the original budget range $0 < \alpha \leq 1/6$, its lower bound is

$$(cN^\beta)^{fg/2} - 1 = c^{fg/2} N^{2\gamma(\alpha)} - 1. \tag{II.210.5}$$

For each fixed α , the positive constant $c^{fg/2}$ and exponent $\gamma(\alpha) > 0$ imply that this is at least $N^{\gamma(\alpha)}$ for sufficiently large N : putting $x = N^{\gamma(\alpha)}$, the inequality $c^{fg/2} x^2 - 1 \geq x$ holds eventually. Thus the clean power in (II.210.4) is consistent with the more local, constant-bearing bound; no new optimal global exponent is claimed. The source theorem additionally covers the full density range up to one, which our original six-budget corollary does not assert.

The older fixed-density quantifier order does not itself give one threshold depending only on a positive lower density bound for all varying actual densities above it. The later uniform theorem in Section II.210.6 supplies that stronger quantifier order. There is a precise integrality issue with the older exact-density quantifier order. For a fixed irrational α , the identity $|A| = \alpha N^2$ is impossible for every positive integer N . For a fixed rational $\alpha = a/b \in (0, 1)$ in lowest terms, integrality forces $b \mid N^2$. With N prime and $b > 1$, this can hold only for a prime dividing the fixed denominator b , and hence for no sufficiently large prime. Thus the pointwise fixed-exact-density statement is eventually vacuous for every fixed $\alpha \in (0, 1)$; $\alpha = 1$ is the remaining case. This does not make the stated theorem false or invalidate its constructive proof. It explains why that older theorem type alone cannot provide a nonvacuous threshold uniform over the actual densities of arbitrary admissible contexts. Our strip contexts have varying actual density for arbitrarily large primes, so the finite-budget obstruction and its compatible re-selection examples do not have this vacuity. This limitation concerns the older API, not the latest uniform theorem. Reconciliation with unchanged catalogue or printed target statements remains a separate question outside this report. In particular, (II.210.4) does not assert a grid for an arbitrary previously selected Stage 13.9 witness or a side

bound in terms of that witness’s actual length. We inspected and byte-verified this source but did not rerun its Lean build or axiom checks. The later source amendment below completes the report’s dated comparison.

[write] **7 October 2026; sources 47–53.** *The same quantifier form at HEAD 06b30ca52.* The repository’s later theorems `section13_complete_square_extraction_below_limit` and `section13_complete_square_extraction_three_halves` (`Proofs13SquareExponentLimit.lean`, lines 63 and 146) reach every exponent below $2^{-385}\alpha^{1856}/2^{10Q_{13}(\alpha)}$, using an improved Stage 13.6 length (`lemma_13_6_improved_length`, `Proofs13ImprovedRowLength.lean`, line 11); they are stated for contexts of exact density α (`S.alpha=alpha`, with $|A| = \alpha N^2$ in `Section13Context`), so the integrality observation above applies to them verbatim: for fixed $0 < \alpha < 1$ they hold vacuously once N exceeds the denominator of α . This is an observation about their quantifiers, not a defect in their proofs, and they are not compared with the uniform theorems.

II.210.6 Uniform actual-density extraction is now supplied upstream

The final source pin for this report is `4f8cd11982e6037060b594a3da4296260d26f715`, checked on 7 October 2026. In `Proofs13UniformSquareExtraction.lean`, lines 15–24, the theorem `section13_complete_square_extraction_uniform_density` has the correct uniform order

$$\forall \rho > 0 \ \exists N_0(\rho) \ \forall N \geq N_0(\rho) \text{ prime } \forall \text{ contexts with } \rho \leq \alpha.$$

Here ρ denotes the source theorem’s lower density bound, distinguished from the coefficient-set density δ used in our repair. It constructs proper equal-length progressions V, W of the same nonzero step and a set $B \subseteq A \cap (V \times W)$ on which the original ϕ is bilinear, with

$$\ell(V) = \ell(W) \geq N^{\gamma(\rho)}, \quad |B| \geq 2^{-137} \rho^{704} \ell(V) \ell(W). \quad (\text{II.210.6})$$

The function γ is the explicit positive function in (II.210.4). Since contexts have $\alpha \leq 1$, the nonempty input range is $0 < \rho \leq 1$; their actual density may vary anywhere in $[\rho, 1]$ as N varies.

This source theorem resolves the basic uniformity limitation of the older exact-density statement. The integrality observation in Section II.210.5 remains a historical distinction between the two APIs, not a missing feature of the latest source. The new result still does not assert the exact grid for arbitrary already retained Stage 13.9 witnesses, so Theorem F.56.1 is unaffected. We credit the declared uniform theorem and its source proof, without claiming to have rerun its Lean compilation or axiom audit. No stronger actual-density output or explicit uniform threshold is added to this report, and the source comparison is frozen at this pin.

[write] **7 October 2026; sources 47–53.** *Explicit thresholds since the pin* (checked at HEAD 06b30ca52). The repository now states the density-uniform extraction with a fully explicit threshold, `section13_complete_square_extraction_explicit` (`Proofs13ExplicitGeometricThreshold.lean`, line 24; threshold `section13GeometricThreshold`, line 12; commit 93d436414), with side $N^{\gamma(\delta)}$ and mass $2^{-137}\delta^{704}$ at the lower density δ ; source 53 gives a closed threshold and the actual-density side $N^{3\gamma(a)}$ (Section II.195). Neither changes Theorem F.56.1.

II.210.7 A restricted sharpness statement

The uniform factor $1/4$ is asymptotically sharp if one insists on the prescribed step and side $M - 1$. Work in an integer $M \times M$ box, with $M \geq 3$, and place one unit of mass at each of its four corners. The ambient density is $4/M^2$. Every contained consecutive $(M - 1) \times (M - 1)$ square contains

exactly one corner and has density $1/(M-1)^2$. The ratio of output to ambient density is

$$\frac{M^2}{4(M-1)^2} \longrightarrow \frac{1}{4}.$$

The same example embeds without wrap in a sufficiently large cyclic group. This proves only the stated fixed-step, prescribed-length sharpness. It does not prove optimality if arbitrary other steps or shorter output progressions are allowed, and it is not a sharpness statement for the square-root side length used in the source wrapper.

II.211 A direct repair for suitably selected witnesses

Source 51, delivered Section 10.

The obstruction identifies information that a lossy retained interface cannot recover for arbitrary witnesses. It does not prevent a better selection from carrying that information. We now prove an existential repair by preserving the short parent already used in the actual Stage 13.6 construction, localizing the coefficient extraction, and taking a final square of side $\ell(U) - 1$.

The geometric arguments below are self-contained. One analytic input is imported from the pinned source and is stated explicitly. We do not reprove that analytic theorem or assert a new Lean verification of the resulting repair.

II.211.1 The imported all-scale coefficient input

Source input. The theorem `extract_bilinear_cell` in `Proofs13CoefficientPartition.lean`, lines 84–127 at commit 2548e7fffc64de1a05eadae81b9247d9b9042c3b, supplies the following statement. Let N be prime, T a proper progression of nonzero step and positive length v , X_0 a nonempty column set, and $J_0 \subseteq T$. Suppose $0 < \delta \leq 1$, the maps A_c, B_c are order-eight Freiman on J_0 , and

$$C_0 \subseteq X_0 \times (J_0 + y), \quad |C_0| \geq \delta |X_0| v, \quad \phi(x, z) = A_c(z - y) + B_c(z - y)x \quad ((x, z) \in C_0).$$

Then there is a proper nonzero-step progression $U \subseteq T$ such that, with

$$e = \delta^2 / 2^{14}, \quad D = \{(x, z) \in C_0 : z - y \in U\}, \quad (\text{II.211.1})$$

one has

$$\ell(U) \geq v^e, \quad |D| \geq \delta |X_0| \ell(U), \quad \phi \text{ is bilinear on } D.$$

There is no lower-scale premise such as $8 \leq v^e$ in this source theorem. Its proof applies the existing all-scale simultaneous affine partition to the two coefficient maps on their one shared domain, then uses weighted selection. For $q \geq 1$, the definition in `Proofs07SimultaneousLinearity.lean`, lines 1075–1077, is `cor711Exponent`(δ, q) = $2^{-14} \delta^2 / q$, giving the displayed one-domain exponent at $q = 1$. In particular, the two maps do not incur two independent exponent losses.

II.211.2 Short-carrier rectification and balanced parents

Lemma II.211.1 (Rectification inside a short progression). *Let P be a proper progression in $\mathbb{F}_N$ of nonzero step d and length n_P , with $2(n_P - 1) < N$. If a proper progression X_1 of length at least two*

has carrier contained in P , then, after possibly reversing its indexing, it has step ad for a positive ordinary integer a satisfying

$$a(\ell(X_1) - 1) \leq n_P - 1. \quad (\text{II.211.2})$$

Reversal leaves its carrier and length unchanged.

Proof. Normalize P by the affine bijection $z \mapsto (z - P.\text{start})/d$. The successive points of X_1 have unique indices $i_j \in \{0, \dots, n_P - 1\}$. Every difference $i_{j+1} - i_j$ is in $[-(n_P - 1), n_P - 1]$ and represents the same residue d_{X_1}/d . Two such integer representatives differ in absolute value by at most $2(n_P - 1) < N$, so they are equal as integers. Write their common value as a_0 ; it is nonzero because the progression is proper and has at least two points. Thus $i_j = i_0 + ja_0$. Reverse the progression if $a_0 < 0$ and put $a = |a_0|$. The difference of its extreme indices proves (II.211.2). $\square$

Lemma II.211.2 (Balanced aligned parent partition). *Let $L \geq 2$, $a \geq 1$, and $a(L - 1) < q_Q$ be integers. The interval $\{0, \dots, q_Q - 1\}$ has a disjoint partition into step- a progressions, each of length between $\lceil L/2 \rceil$ and L .*

Proof. Split the interval into its a residue chains. Since $q_Q \geq a(L - 1) + 1$, every chain has length $r_a \geq L - 1$, and all chains are nonempty. For one chain of length r_a , put $c_a = \lceil r_a/L \rceil$ and split it into c_a consecutive blocks whose lengths differ by at most one. Their lengths are $\lfloor r_a/c_a \rfloor$ or $\lceil r_a/c_a \rceil$, and they are at most L . If $c_a = 1$, the chain length lies in $[L - 1, L]$ and hence is at least $\lceil L/2 \rceil$. If $c_a \geq 2$, then

$$r_a \geq (c_a - 1)L + 1, \quad \frac{r_a}{c_a} \geq L - \frac{L - 1}{c_a} \geq \frac{L + 1}{2}.$$

Thus the smallest block has length at least $\lfloor (L + 1)/2 \rfloor = \lceil L/2 \rceil$. The blocks partition each chain, so together they partition the original interval. $\square$

The lower bound is optimal for this class of partitions: a chain of length $L + 1$ needs at least two blocks of length at most L , and one block then has length at most $\lfloor (L + 1)/2 \rfloor = \lceil L/2 \rceil$. This is a statement about the parent partition, separate from the density sharpness in Section II.210.

II.211.3 The direct geometric and analytic assembly

Theorem II.211.3 (Short-parent square extraction). *Let N be prime. Let R, Q, S be proper progressions in $\mathbb{F}_N$, with all three steps nonzero, $d_R = d_Q$, and $S \subseteq R$. Write*

$$n_R = \ell(R), \quad q_Q = \ell(Q), \quad L = \ell(S).$$

Assume $L \geq 3$, $n_R \leq q_Q$, and $2(n_R - 1) < N$. Let $0 < \delta \leq 1$, $J_0 \subseteq Q$, and suppose

$$\begin{aligned} C &\subseteq S \times (J_0 + y), & |C| &\geq \delta L q_Q, \\ \phi(x, z) &= A_c(z - y) + B_c(z - y)x \quad ((x, z) \in C). \end{aligned} \quad (\text{II.211.3})$$

where the coefficient pair (A_c, B_c) is order-eight Freiman on J_0 . Using the source input (II.211.1), there exist a newly selected proper nonzero-step progression $U \subseteq Q$, proper common-step progressions $V \subseteq S$, $W \subseteq U + y$ of the same nonzero step and the same length k , and a set $E \subseteq C \cap (V \times W)$, such that

$$\begin{aligned} k &= \ell(U) - 1 \geq \lceil L/2 \rceil^{\delta^2/2^{14}} - 1 \geq L^{\delta^2/2^{15}} - 1, \\ |E| &\geq (\delta/4)k^2, \quad \phi \text{ is bilinear on } E. \end{aligned} \quad (\text{II.211.4})$$

Proof. Apply Lemma II.211.1 to $S \subseteq R$. Represent the same carrier S , reversing its indexing if necessary, with

$$d_S = ad_Q, \quad a \geq 1, \quad a(L-1) \leq n_R - 1 < q_Q.$$

Transport Lemma II.211.2 through the injective indexing of Q . This partitions Q into proper parents T_i of step d_S and lengths

$$h := \lceil L/2 \rceil \leq v_i := \ell(T_i) \leq L.$$

Every parent is short enough to rectify its subprogressions, since $2(v_i - 1) \leq 2(L-1) \leq 2(n_R - 1) < N$.

Let $C_i = \{(x, z) \in C : z - y \in T_i\}$. This is a disjoint partition of C , and $\sum_i v_i = q_Q$. Weighted averaging therefore gives one parent $T = T_i$, of length v , with

$$|C_i| \geq \delta L v. \quad (\text{II.211.5})$$

This uses actual pair mass, not merely the number of retained heights. Put $J_T = J_0 \cap T$ and $C_T = C_i$. Its support is $S \times (J_T + y)$; the original row formula and the Freiman coefficient pair restrict to J_T . Projection of the pair onto either coordinate gives the two scalar Freiman maps. Also $|C_T| \leq L|J_T|$, so (II.211.5) implies $|J_T| \geq \delta v$, as needed within the analytic source proof.

Invoke (II.211.1) with column set S , parent T , domain J_T , set C_T , and the same $y, \phi, A_c, B_c, \delta$. It supplies $U \subseteq T$ with

$$M := \ell(U) \geq v^e, \quad e = \delta^2/2^{14}, \quad D = \{(x, z) \in C_T : z - y \in U\}, \quad |D| \geq \delta L M,$$

and bilinearity on D . No density is lost in this localization. Because $U \subseteq T$ and C was supported over J_0 , the same set is exactly

$$D = \{(x, z) \in C : z - y \in U \cap J_0\}.$$

Since $L \geq 3$, we have $v \geq 2$ and $v^e > 1$, so the integer M is at least two. Apply Lemma II.211.1 again to $U \subseteq T$, reversing U if necessary. Its carrier and D stay unchanged, while its step becomes

$$d_U = bd_S, \quad b \geq 1, \quad b(M-1) \leq v-1 \leq L-1.$$

Thus Proposition II.210.1 applies to S and $U + y$ with the stronger choice $k = M - 1$. It gives the asserted V, W, E , density, containments, and bilinearity. Finally,

$$k = M - 1 \geq v^e - 1 \geq \lceil L/2 \rceil^e - 1 \geq L^{e/2} - 1.$$

The last inequality follows from $\lceil L/2 \rceil^2 \geq L$ for every integer $L \geq 3$. □

For $L = 1$ or 2 , positive mass makes C nonempty. One point of C gives a proper singleton square at any chosen nonzero step, carrying density one and a constant bilinear value. It meets the weaker effective target $L^{\delta^2/2^{15}} - 1$. These cases are not needed under the actual width budget.

II.211.4 The original phase and the height offset are preserved

All selections above restrict original carriers or reverse their indexing. The phase is always the original ϕ , evaluated in the original coordinates. To see the height offset explicitly, the analytic source proof produces affine row coefficients on $U \cap J_0$ of the form

$$A_c(h) = a_1 h + a_0, \quad B_c(h) = b_1 h + b_0.$$

For an actual point $(x, z) \in D$, substitute $h = z - y$ to obtain

$$\phi(x, z) = (a_0 - a_1 y) + (b_0 - b_1 y)x + a_1 z + b_1 xz. \quad (\text{II.211.6})$$

This is a bilinear formula in the original variables, with both intercept corrections included. Reversing S or U does not replace x or h by a new coordinate: it merely reorders the same carrier. Consequently the ambient functions and the formula (II.211.6) are unchanged. In the actual stage chain, $E \subseteq C \subseteq B \subseteq A$, so the original-domain containment is also preserved.

II.211.5 Recovering the short parent from the actual budget construction

The hypotheses of Theorem II.211.3 can be obtained for suitably selected outputs under the existing six budgets. This requires following the actual construction, rather than taking an arbitrary witness from its lossy existential conclusion.

The source `common_step_row_selection` in `Proofs13CommonStepSelection.lean`, lines 92–132, selects R with

$$\ell(R) = m \quad \text{or} \quad \ell(R) = m + 1. \quad (\text{II.211.7})$$

The proof of `lemma_13_6_of_short_difference_rules` receives this fact as its local length alternative and constructs F from that same R . The subsequent Bohr-model and initial-stage wrappers do not alter R , but their stated conclusions retain only $\ell(R) \leq \ell(Q)$. Thus (II.211.7) can be preserved by using those same witness constructions and exporting the length alternative. It cannot be inferred for every arbitrary F satisfying the wrapper's weaker conclusion.

Put $A_0 = cN^\beta$, with c, β, f, W as in (II.203.4), and let $L = \ell(S)$ for the Stage 13.7 output. The existing lower-length and width budgets, followed by the Stage 13.7 length bound, give

$$L \geq (\ell(R))^f \geq A_0^f \geq W = \delta^{-1} \geq 2^{135}, \quad \delta = 2^{-135} \alpha^{704}.$$

Properness and $S \subseteq R$ give $L \leq \ell(R) \leq m + 1$, so $m \geq 3$. Consequently

$$2(\ell(R) - 1) \leq 2m < m^2 \leq N. \quad (\text{II.211.8})$$

The steps of R, Q agree and are nonzero by Stage 13.6, while $\ell(R) \leq m + 1 \leq \ell(Q)$. The properness and nonzero step of S are explicit Stage 13.7 conditions. The Stage 13.8 support, mass, row formula and pair Freiman property give (II.211.3). Reversing the indexing of S preserves these stages: their support and row statements concern its carrier and the actual first coordinate.

Corollary II.211.4 (Existential square extraction under the existing budgets). *For every prime N , context, Stage 13.4 and 13.5 data, and integer m satisfying all the hypotheses of the pinned budget theorem, its actual constructions can be used to select F, G, H satisfying the exact Stage 13.6–13.8 predicates and (II.211.7), and then a proper same-step bilinear square with retained set in $C \subseteq A$, density at least $2^{-137} \alpha^{704}$, and side k satisfying*

$$\begin{aligned} k &\geq \lceil L/2 \rceil^{2^{-284} \alpha^{1408}} - 1 \geq L^{2^{-285} \alpha^{1408}} - 1 \\ &\geq (cN^\beta)^{2^{-385} \alpha^{1856}} - 1. \end{aligned} \quad (\text{II.211.9})$$

This is a mathematical consequence using the imported analytic source input; no new formal implementation of this corollary is asserted.

Proof. The preceding argument supplies every hypothesis of Theorem II.211.3. Here $e = \delta^2/2^{14} = 2^{-284}\alpha^{1408}$ and $\delta/4 = 2^{-137}\alpha^{704}$. The last line of (II.211.9) follows from $L \geq (cN^\beta)^f$ and $f(e/2) = 2^{-385}\alpha^{1856}$. $\square$

[write] **7 October 2026; sources 47–53.** *Exponents* (checked at the write). With $e = 2^{-284}\alpha^{1408}$ and $f = 2^{-100}\alpha^{448}$, the last line is $(cN^\beta)^{2^{-385}\alpha^{1856}} - 1 = c^{f/2}N^{2\gamma(\alpha)} - 1$, $\gamma(\alpha) = \beta f e/4 = 2^{-386}\alpha^{1856}2^{-13Q_{13}(\alpha)}$, the formal **section13SquareExponent**: exponent 2γ . Source 53's Proposition II.201.4 is the same assembly, and keeps the first bound $\lceil L/2 \rceil^e - 1 \geq (L/2)^e - 1$ instead of passing to $L^{e/2}$; that is its exponent $4\gamma(a)$, with a uniform explicit threshold (Theorem F.58.1).

The effective density and side bounds are the ones obtained by composing the recorded lower length in Stage 13.9 with the intended square-root extraction in Corollary 13.10. The new localized U is only proved to have length at least $\lceil L/2 \rceil^e$, not the unchanged Stage 13.9 target L^e . Nor is the new square compared to the actual length of an independently selected old Stage 13.9 witness. The route goes directly from newly chosen Stage 13.8 data to a square and bypasses that intermediate encoding. It therefore coexists with the bad retained witnesses in Theorem F.56.1.

II.212 The precise retained interface

Source 52, delivered Section 2.

We translate the source into mathematical notation before choosing witnesses. All finite counts are unnormalized, and all real exponent laws below have positive bases.

II.212.1 Progressions and ambient affine maps

A modular progression $P = (a, d, L)$ has carrier

$$[P] = \{a + jd : 0 \leq j < L\} \subseteq \mathbb{Z}/N\mathbb{Z}, \quad \ell P = L.$$

It is proper when $|[P]| = L$. The definition concerns injectivity of the parameterization; it does not forbid a carrier that crosses the conventional residue cut. Nonzero step is a separate condition. The predicate **LinearOn** on $U \subseteq \mathbb{Z}/N\mathbb{Z}$ means

$$\exists a, b \in \mathbb{Z}/N\mathbb{Z} \quad \forall x \in U, \quad f(x) = ax + b. \quad (\text{II.212.1})$$

It does not mean merely that $f(a_0 + jd_0)$ is affine as a function of the integer index j .

II.212.2 Context and arrangement normalization

An order-eight Freiman homomorphism on V preserves equality between sums of eight elements of V , with repetitions allowed. A context comprises $A \subseteq (\mathbb{Z}/N\mathbb{Z})^2$, a total function $\phi : (\mathbb{Z}/N\mathbb{Z})^2 \rightarrow \mathbb{Z}/N\mathbb{Z}$, and reals α, η satisfying

$$0 < \alpha \leq 1, \quad |A| = \alpha N^2, \quad \eta = 2^{-44},$$

the separate order-eight Freiman property on all horizontal and vertical sections of A , and the required mostly-respected inequality.

To specify the latter, an eight-arrangement at fixed height h consists of sixteen x_i and sixteen y_i with

$$\sum_{i=1}^8 x_i = \sum_{i=9}^{16} x_i, \quad (x_i, y_i), (x_i, y_i + h) \in A \quad (1 \leq i \leq 16).$$

It is respected if the first eight and last eight sums of $\phi(x_i, y_i + h) - \phi(x_i, y_i)$ agree. Write $C(h)$ for the number of contained arrangements and $C_\phi(h)$ for the number of respected contained arrangements. The context requires

$$(1 - \eta) \sum_h C(h) \leq \sum_h C_\phi(h).$$

A height is good if $(1 - 2\eta)C(h) \leq C_\phi(h)$, and strong if it is good and

$$C(h) \geq \alpha^{32} N^{31} / 16. \quad (\text{II.212.2})$$

For records D, E , define

$$\mathcal{H} = ([E.Q] \cap D.H) \cap \{h : h \text{ is strong}\}.$$

This is exactly `criticalHeights`.

II.212.3 Stage 13.6 in full

Write

$$\begin{aligned} \delta &= 2^{-43} \alpha^{224}, & \lambda &= 2^{-48} \alpha^{256}, & \varepsilon &= 2^{-100} \alpha^{448}, \\ K_{13} &= 2^{114} \alpha^{-320}, & Q_{13} &= 2^{220} \alpha^{-221}, & \zeta_{13} &= 2^{-228 K_{13}} \alpha^{576 K_{13}}, \\ c &= \zeta_{13} / 2, & \beta &= 2^{-13 Q_{13}}. \end{aligned}$$

For $h \in \mathbb{Z}/N\mathbb{Z}$, put

$$\mathcal{E}_h = \{(x, y) : (x, y), (x, y + h) \in A\}, \quad Y_h^R = \{(x, y) \in F.Y(h) : x \in [F.R]\}.$$

The full `IsStage136Data` predicate is the conjunction of:

136.1 $F.R$ and $E.Q$ have the same nonzero step, and $F.R$ is proper;

136.2 $cN^\beta \leq \ell(F.R)$;

136.3 for each $h \in \mathcal{H}$, $F.Y(h) \subseteq \mathcal{E}_h$, $\delta N^2 \leq |F.Y(h)|$, and the vertical difference of ϕ is of the form $a_h x + b_h$ on Y_h^R ;

136.4 $\delta \ell(F.R) N |\mathcal{H}| \leq \sum_{h \in \mathcal{H}} |Y_h^R|$;

136.5 $\lambda \ell(E.Q) \ell(F.R) N \leq \sum_{h \in \mathcal{H}} |Y_h^R|$.

In particular there is no conjunct $\ell(F.R) \leq \ell(E.Q)$ and no requirement that $F.R$ lie in $E.Q$.

II.212.4 The two Stage 13.7 conclusions

A candidate output G consists of $y_0 \in \mathbb{Z}/N\mathbb{Z}$, a progression S , and $B \subseteq (\mathbb{Z}/N\mathbb{Z})^2$. Write $L = \ell S$. The weak predicate comprises:

137.1 S has nonzero step, is proper, and $[S] \subseteq [F.R]$;

137.2 $(\ell(F.R))^\varepsilon \leq L$;

137.3 $B \subseteq [S] \times (y_0 + \mathcal{H})$;

137.4 $\delta L |\mathcal{H}| \leq |B|$;

137.5 $\lambda \ell(E.Q)L \leq |B|$;

137.6 for every $h \in \mathcal{H}$, the map $x \mapsto \phi(x, y_0 + h)$ is ambient affine on $U_h = \{x \in [S] : (x, y_0 + h) \in B\}$.

The full predicate is precisely this weak predicate together with $B \subseteq A$. We will contradict only conditions 137.1–137.4 and 137.6. Dropping the second mass lower bound or the original-domain requirement cannot repair the examples.

II.212.5 The genuine initial-stage predicates

For the strengthening we also need the complete initial-stage interface. Define

$$\theta_1 = 2^{-1882}\theta^{10477}, \quad Q_{134}(\theta) = 2^{1882}\theta^{-10479}.$$

The record D stores q, m_D, P, H_{134} and q coefficient pairs (a_i, b_i) . The Stage 13.4 conditions are:

134.1 $q > 0$, P proper with nonzero step, and $H_{134} \subseteq [P]$;

134.2 $q \leq Q_{134}(\theta)$;

134.3 $m_D \leq \theta_1 N^{\theta_1^2/(16q)}/(64\pi) < m_D + 1$;

134.4 $\ell P = m_D$ or $\ell P + 1 = m_D$;

134.5 $\alpha^{32} N^{31} \ell P / 8 \leq \sum_{h \in H_{134}, h \text{ good}} C(h)$;

134.6 for $h \in H_{134}$ and every r with $|\widehat{f}_h(r)| \geq \theta N^2$, there is i with $r = a_i h + b_i$.

Here $f_h(x) = |\{y : (x, y) \in \mathcal{E}_h\}|$, and the source transform is unnormalized:

$$\widehat{f}(r) = \sum_{x \in \mathbb{Z}/N\mathbb{Z}} f(x) \exp(\pm 2\pi i r x / N).$$

The sign convention has no effect on the bounds in this report. The separate theorem `lemma_13_4` requires primality and $0 < \theta \leq \alpha^{32}/4$. Our construction verifies its conclusion predicate directly and satisfies this numerical threshold condition, but does not assert primality.

The Stage 13.5 record stores Q . Its predicate requires:

135.1 Q proper with nonzero step and $[Q] \subseteq [P]$;

$$135.2 \quad (\ell P)^{2^{-12q}}/2 \leq \ell Q;$$

$$135.3 \quad \text{for every } i \text{ and } h \in [Q], \|(a_i h + b_i)d_Q\|_N \leq (\ell P)^{-2^{-11q}}N;$$

$$135.4 \quad \alpha^{32}\ell Q/20 \leq |\mathcal{H}|.$$

The quantity $\|r\|_N$ is the least absolute value of an integer representative of r . This spells out every extra premise used later.

II.213 The common dyadic strip

Source 52, delivered Section 3.

For this section take any integer $k \geq 5$ and set

$$N = 2^k, \quad m = 2^{k-4}, \quad T = \{2j \bmod N : 0 \leq j < m\}, \quad A = T \times \mathbb{Z}/N\mathbb{Z}.$$

Let $P_0 = (0, 2, m)$ and define the total function

$$\phi(x, y) = \left\lfloor \frac{x.\text{val}}{2} \right\rfloor \pmod{N}, \tag{II.213.1}$$

where $x.\text{val} \in \{0, \dots, N-1\}$ is the canonical representative. In particular $\phi(2j, y) = j$ for $0 \leq j < m$, and ϕ is independent of y everywhere.

Lemma II.213.1 (The context is genuine). *With $\alpha = 1/16$ and $\eta = 2^{-44}$, these data satisfy every field of Section13Context. Moreover every height is strong.*

Proof. The representatives $2j$ are distinct and less than N , so $|T| = m$ and P_0 is proper with nonzero step. Hence $|A| = mN = \alpha N^2$, with $0 < \alpha \leq 1$.

Every nonempty vertical section is the full group, and the corresponding map is constant. Empty vertical sections give a vacuous Freiman condition. For a horizontal section, take indices $j_1, \dots, j_8, j'_1, \dots, j'_8 \in \{0, \dots, m-1\}$ whose doubled sums agree modulo N . Put

$$\Delta = \sum_{i=1}^8 j_i - \sum_{i=1}^8 j'_i \in \mathbb{Z}.$$

The domain equality implies $N \mid 2\Delta$, hence $(N/2) \mid \Delta$. But

$$|\Delta| \leq 8(m-1) = N/2 - 8 < N/2.$$

Thus $\Delta = 0$, and the image sums agree even as integers. This proves the order-eight condition with repeated entries allowed.

All vertical differences of ϕ vanish. Every contained arrangement is consequently respected, globally and at each height. Thus $C_\phi(h) = C(h)$, which implies the mostly-respected and good-height inequalities because $\eta > 0$.

For the strong-height lower bound, let $r(s)$ be the number of ordered eight-tuples in T with sum s , and put $E_8(T) = \sum_s r(s)^2$. There are m^8 such tuples in total. Cauchy–Schwarz over the N possible sums gives

$$E_8(T) \geq m^{16}/N. \tag{II.213.2}$$

For each fixed height h , an additive sixteen-tuple of first coordinates permits exactly N^{16} independent choices of the second coordinates: both endpoints of every resulting edge belong to A . Therefore

$$C(h) = C_\phi(h) = N^{16}E_8(T) \geq m^{16}N^{15} = \alpha^{16}N^{31} \geq \alpha^{32}N^{31}/16. \quad (\text{II.213.3})$$

The height is fixed in this count, so there is no extra factor N . The sixteen second coordinates are independent, so none of those sixteen factors is missing. This proves (II.212.2) at every height. $\square$

Lemma II.213.2 (Uniform edge data). *For every height, $\mathcal{E}_h = A$, $f_h = N\mathbf{1}_T$, and the vertical difference of ϕ is zero. If a progression R contains T and $F.Y(h) = A$, then $Y_h^R = A$ and its affine difference coefficients can be taken to be zero.*

Proof. The two edge endpoints lie in A exactly when their common first coordinate is in T . All remaining assertions follow directly. $\square$

At $\alpha = 2^{-4}$ the constants become

$$\begin{aligned} \delta &= 2^{-939}, & \lambda &= 2^{-1072}, & \varepsilon &= 2^{-1892}, \\ K_{13} &= 2^{1394}, & Q_{13} &= 2^{9 \cdot 2^{20}}, & c &= 2^{-2532K_{13}-1}. \end{aligned} \quad (\text{II.213.4})$$

Indeed $43 + 4 \cdot 224 = 939$, $48 + 4 \cdot 256 = 1072$, $100 + 4 \cdot 448 = 1892$, $114 + 4 \cdot 320 = 1394$, and $228 + 4 \cdot 576 = 2532$. Also $2^{20} + 4 \cdot 2^{21} = 9 \cdot 2^{20}$. In particular

$$0 < c < \alpha, \quad 0 < \beta < 1, \quad 0 < \delta < \alpha < 1, \quad 0 < \lambda < 1. \quad (\text{II.213.5})$$

For every $N \geq 1$, it follows that

$$cN^\beta \leq cN < \alpha N = m. \quad (\text{II.213.6})$$

This settles the Stage 13.6 length condition without evaluating K_{13} , Q_{13} , or N numerically.

II.214 The exact step two counterexample

Source 52, delivered Section 4.

We now prove Theorem F.57.1. Take its specified k and the common context above. Set

$$\begin{aligned} D.q &= 1, & D.m &= m, & D.P &= P_0, & D.H &= T, & D.a(0) &= D.b(0) = 0, \\ E.Q &= P_0, & F.R &= P_0, & F.Y(h) &= A & (h \in \mathbb{Z}/N\mathbb{Z}). \end{aligned}$$

All heights are strong, so $\mathcal{H} = T$ and $|\mathcal{H}| = m > 0$.

II.214.1 Every required premise

The common step is the nonzero residue 2, and $F.R$ is proper. Equation (II.213.6) proves the progression length bound. Lemma II.213.2 gives the exact edge-domain and restricted-edge identities, and the zero affine vertical difference. Edge density follows from $|A| = \alpha N^2 \geq \delta N^2$.

The sum in the last two Stage 13.6 conjuncts is exactly

$$\sum_{h \in T} |Y_h^R| = m|A| = m^2N.$$

The required lower bounds are respectively $\delta m^2 N$ and $\lambda m^2 N$, so both hold by (II.213.5). This verifies the whole predicate, rather than appealing to the existence theorem of Stage 13.6.

The extra Stage 13.5 predicate also holds. Its geometric assertions are immediate from $Q = P = P_0$; its length condition follows from $m \geq 1$ and $0 < 2^{-12} \leq 1$; all coefficient expressions vanish; and its critical-height condition is $\alpha^{32} m / 20 \leq m$. No Stage 13.4 certificate is inferred from the record name or these checks.

II.214.2 The affine obstruction

Lemma II.214.1 (At most one point per row). *For any $a, b, y \in \mathbb{Z}/N\mathbb{Z}$, the equation $\phi(x, y) = ax + b$ holds at at most one point $x \in T$.*

Proof. If it holds at $2j$ and $2j'$ with $0 \leq j, j' < m$, subtraction gives

$$(2a - 1)(j - j') = 0 \quad \text{in } \mathbb{Z}/2^k\mathbb{Z}.$$

The residue $2a - 1$ is odd, so it is coprime to 2^k and is a unit. Hence $j = j'$ modulo N . As $j, j' \in [0, m)$ and $m < N$, they are equal as integers. Thus the two points coincide. $\square$

This argument uses the prime-power structure. For an arbitrary even composite modulus an odd residue need not be a unit; the mere fact that step 2 is a nonunit would not justify the same proof.

II.214.3 The incompatible output bounds

Suppose a weak Stage 13.7 output existed. Its row supports U_h lie in T , so Lemma II.214.1 and the affine row condition give $|U_h| \leq 1$. Translation $h \mapsto y_0 + h$ is injective. The support condition therefore partitions B into exactly the possible rows indexed by T , and

$$|B| = \sum_{h \in T} |U_h| \leq m.$$

The first mass lower bound yields

$$\delta L m \leq |B| \leq m, \quad L \leq \delta^{-1} = 2^{939}.$$

The output length condition yields the strict opposite:

$$\begin{aligned} L &\geq m^\varepsilon = \left(2^{k-4}\right)^{2^{-1892}} = 2^{(k-4)2^{-1892}} \\ &= 2^{939+2^{-1892}} > 2^{939}. \end{aligned}$$

This contradiction proves Theorem F.57.1. In this particular construction, the support already implies $B \subseteq T \times \mathbb{Z}/N\mathbb{Z} = A$, so the missing original-domain conjunct is not the source of failure.

II.215 Certifying both genuine initial stages

Source 52, delivered Section 5.

This section gives the step-two version of the second construction, to isolate the additional Fourier and floor checks from the later unit-step modification. Keep $\alpha = 1/16$ and set

$$\begin{aligned} \theta &= 2^{-130} = \alpha^{32}/4, & t &= 1882 + 130 \cdot 10477 = 1\,363\,892, \\ K_0 &= 2^{128}, & q &= 2(2K_0 + 1) = 2^{130} + 2, \\ U &= 16q\,2^{2t}, & u &= 1/U, \\ k &= (t + 10)U, & N &= 2^k, \quad m = N/16. \end{aligned} \tag{II.215.1}$$

All these numbers are exact finite integers or positive rationals as indicated. We never need their decimal expansions. The same context from Section II.213 is valid, as $k \geq 5$.

Let $P_1 = (0, 2, 5)$, set $D.P = E.Q = P_1$, $D.m = 5$, and $D.H = [P_1]$. Index the q pairs (ν, s) with $\nu \in \{0, 1\}$ and $-K_0 \leq s \leq K_0$ by any bijection with $\text{Fin } q$, and define

$$a_{\nu,s} = 0, \quad b_{\nu,s} = \nu N/2 + s \pmod{N}. \tag{II.215.2}$$

Keep $F.R = P_0$ and $F.Y(h) = A$. The ordering of the coefficient graphs has no mathematical role.

II.215.1 A uniform cover of the large Fourier spectrum

Lemma II.215.1 (The two alias families). *For every height h and frequency r ,*

$$|\widehat{f}_h(r)| \geq \theta N^2 \implies r = \nu N/2 + s \pmod{N} \text{ for some } \nu \in \{0, 1\}, |s| \leq K_0.$$

Every graph value in (II.215.2) also satisfies $\|2b_{\nu,s}\|_N \leq 2^{129}$.

Proof. By Lemma II.213.2, $f_h = N\mathbf{1}_T$. Thus, up to complex conjugation,

$$|\widehat{f}_h(r)| = N \left| \sum_{j=0}^{m-1} \exp(2\pi i \, 2rj/N) \right|.$$

Set $d = \|2r\|_N$, so $0 \leq d \leq N/2$. If $d > 0$, the finite geometric-series formula gives

$$\left| \sum_{j=0}^{m-1} \exp(2\pi i \, 2rj/N) \right| \leq \frac{1}{\sin(\pi d/N)} \leq \frac{N}{2d}.$$

For the last inequality, $\sin(\pi x) \geq 2x$ on $[0, 1/2]$, as follows from the concavity of sine there and its endpoint values. The numerator of the geometric-series quotient has modulus at most 2, and its denominator has modulus $2 \sin(\pi d/N)$. Hence a large coefficient forces

$$\theta N \leq N/(2d), \quad d \leq 1/(2\theta) = 2^{129}.$$

When $d = 0$ the latter bound is automatic; there is no division by zero.

A least-absolute integer representative of $2r$ is even, because the modulus N is even. Write it as $2s$. Then $|s| \leq 2^{128} = K_0$ and $2(r-s) = 0$ in $\mathbb{Z}/N\mathbb{Z}$. The kernel of multiplication by 2 is exactly $\{0, N/2\}$, which gives the two listed aliases. Finally $2b_{\nu,s} = 2s$ modulo N , so $\|2b_{\nu,s}\|_N \leq 2|s| \leq 2^{129}$. $\square$

Both alias families are needed. Replacing the spectral cover by only the interval around frequency zero would omit the frequencies near $N/2$.

II.215.2 Every Stage 13.4 conjunct

The progression P_1 has distinct representatives $0, 2, 4, 6, 8$, and hence is proper with nonzero step. Its carrier equals $D.H$, and $q > 0$. The auxiliary constants simplify exactly to

$$\theta_1 = 2^{-t}, \quad Q_{134}(\theta) = 2^{t+260}, \quad \theta_1^2/(16q) = u.$$

Since $q = 2^{130} + 2 \leq 2^{131} \leq 2^{t+260}$, the graph-count bound holds.

The floor expression is particularly simple:

$$\frac{\theta_1}{64\pi} N^u = \frac{2^{-t}}{64\pi} \left(2^{(t+10)U}\right)^{1/U} = \frac{16}{\pi}. \quad (\text{II.215.3})$$

The elementary bounds $3 < \pi < 22/7$ imply

$$5 < 56/11 < 16/\pi < 16/3 < 6.$$

Thus its floor is $5 = D.m$, and $\ell P_1 = D.m$ satisfies the allowed length alternative.

All five selected heights are good, and $C(h) = C(0) \geq \alpha^{16} N^{31}$. Therefore

$$\sum_{h \in D.H, h \text{ good}} C(h) = 5C(0) \geq 5\alpha^{16} N^{31} \geq \alpha^{32} N^{31} \ell(P_1)/8.$$

The final spectral-cover conjunct is Lemma II.215.1. This checks every Stage 13.4 condition directly at an allowed θ .

II.215.3 Every Stage 13.5 conjunct

Here $Q = P_1$. Properness, nonzero step, and containment are immediate. Since $0 < 2^{-12q} \leq 1$,

$$5^{2^{-12q}}/2 \leq 5/2 \leq 5 = \ell Q.$$

Write $v = 2^{-11q} \in (0, 1]$. For every graph value and every height, Lemma II.215.1 gives

$$\|(a_{\nu,s}h + b_{\nu,s})d_Q\|_N = \|2b_{\nu,s}\|_N \leq 2^{129}.$$

On the other hand, $5^{-v}N \geq N/5 > 2^{129}$ because $k \geq 132$ and $2^{132}/5 > 2^{129}$. This is the required coefficient bound. All heights are strong, so $\mathcal{H} = [P_1]$ and $|\mathcal{H}| = 5$, which also proves $\alpha^{32}\ell Q/20 \leq |\mathcal{H}|$.

II.215.4 Stage 13.6 and the final contradiction

All common edge and length verifications from Section II.214 still apply. The new aggregate edge count is $5mN$, and the required lower bounds are $5\delta mN$ and $5\lambda mN$. Thus Stage 13.6 holds.

Each possible output row has at most one point by Lemma II.214.1. Now $|B| \leq 5$ and $5\delta L \leq |B|$, again giving $L \leq 2^{939}$. To check that the lower length bound is strictly larger, it suffices to observe

$$k = (t+10)16q2^{2t} \geq 2^{2t+4} \geq 2^{1903},$$

$$2^{1903} - 4 - 939 \cdot 2^{1892} = 1109 \cdot 2^{1892} - 4 > 0.$$

Thus $k - 4 > 939 \cdot 2^{1892}$ and $m^\varepsilon > 2^{939}$. No weak output exists.

The progression-length mismatch is explicit: $\ell R = m$ while $\ell Q = 5$. Neither the floor calculation nor the spectral cover supplies a missing upper bound on R . This particular construction does not satisfy that extra bound. Section II.217 supplies a different nonunit-step family that does, together with all six numerical budgets.

II.216 The obstruction persists with unit input step

Source 52, delivered Section 6.

A tempting diagnosis is that the failure comes only from the nonunit step 2 of R . The same context also disproves that proposed repair. Enlarge the input carrier to the ordinary interval

$$R_1 = (0, 1, 2m), \quad [R_1] = \{0, 1, \dots, 2m - 1\}.$$

It contains T and is proper because $2m = N/8 < N$. Its step is a unit. We retain the total extension (II.213.1); its values on odd residues matter for the weak output predicate.

Lemma II.216.1 (At most two points per row). *For any $a, b, y \in \mathbb{Z}/N\mathbb{Z}$, the equation $\phi(x, y) = ax + b$ holds at at most two points of $[R_1]$, and at at most one of those points in T .*

Proof. Every $x \in [R_1]$ has a unique expression $x = 2j + r$ with $r \in \{0, 1\}$ and $0 \leq j < m$, and $\phi(x, y) = j$. Two agreements in the same parity class imply

$$j = a(2j + r) + b, \quad j' = a(2j' + r) + b, \quad (2a - 1)(j - j') = 0.$$

The odd multiplier is a unit modulo 2^k , so $j = j'$. There is at most one agreement per parity. The even class is exactly T . $\square$

II.216.1 A unit-step example with equal input lengths

Set $D.P = E.Q = F.R = R_1$, $D.q = 1$, $D.m = 2m$, $D.H = [R_1]$, and $D.a = D.b = 0$. All heights are strong, hence $\mathcal{H} = [R_1]$ with cardinality $2m$. The Stage 13.5 verification is the earlier zero-coefficient argument with m replaced by $2m$.

Take $F.Y(h) = A$. Its restriction to R_1 remains A . The Stage 13.6 length bound holds because $cN^\beta < m < 2m$. The edge density and zero-difference assertions are unchanged. For either this example or any choice $|\mathcal{H}| = \ell Q = H > 0$, the exact restricted-edge sum is HmN , whereas the required bounds are

$$\delta(2m)NH = 2\delta HmN, \quad \lambda H(2m)N = 2\lambda HmN.$$

Both hold because $2\delta \leq 1$ and $2\lambda \leq 1$.

Any weak output has at most two points in each possible row, so

$$\delta LH \leq |B| \leq 2H, \quad L \leq 2/\delta = 2^{940}.$$

Choose now the explicit exponent

$$k = 940 \cdot 2^{1892} + 4. \tag{II.216.1}$$

Then $2m = 2^{k-3}$, and the lower length bound is

$$L \geq (2m)^\varepsilon = 2^{940+2^{-1892}} > 2^{940}.$$

This gives the unit-step version with $R = Q$. It verifies Stage 13.5, but does not assert the Stage 13.4 predicate for its zero-coefficient record.

II.216.2 A unit-step example with genuine initial stages

Use instead the large parameters (II.215.1) and the same spectral graphs (II.215.2). Replace P_1 by $P'_1 = (0, 1, 5)$, and set

$$D.P = E.Q = P'_1, \quad D.H = [P'_1], \quad D.m = 5, \quad F.R = R_1.$$

All input progressions now have the unit step 1. The Fourier function, coefficient cover, graph-count bound and exact floor expression are unchanged. The five selected heights are still good, so every Stage 13.4 conjunct holds exactly as in Section II.215.

Only the Stage 13.5 coefficient estimate needs a different argument. Every residue has centered absolute value at most $N/2$. Since $q \geq 1$,

$$0 < v = 2^{-11q} \leq 1/4, \quad 5^v \leq 5^{1/4} \leq 2, \quad 5^{-v} \geq 1/2.$$

Consequently

$$\|(a_{\nu,s}h + b_{\nu,s})d_Q\|_N = \|b_{\nu,s}\|_N \leq N/2 \leq 5^{-v}N.$$

The rest of Stage 13.5 is unchanged, and $\mathcal{H} = [P'_1]$ has cardinality five. The Stage 13.6 mass calculation just made applies with $H = 5$.

A weak output would again require $L \leq 2^{940}$. But the same large k obeys

$$k \geq 2^{1903}, \quad 2^{1903} - 3 - 940 \cdot 2^{1892} = 1108 \cdot 2^{1892} - 3 > 0.$$

Hence $(2m)^\varepsilon > 2^{940}$, a contradiction. This proves the unit-step assertion of Theorem F.57.2. It has $\ell R = 2m > 5 = \ell Q$.

II.216.3 Exactly which combinations have been ruled out

Construction	Genuine 13.4	Genuine 13.5	Unit step	$\ell R \leq \ell Q$
Section II.214	not asserted	yes	no	yes
Section II.215	yes	yes	no	no
Unit step with $R = Q$	not asserted	yes	yes	yes
Unit step with initial stages	yes	yes	yes	no
Section II.217	yes	yes	no	yes

Every row has a genuine context and Stage 13.6 predicate, and every row fails even the weak Stage 13.7 conclusion. An unasserted certificate in the table must not be silently supplied. In particular the table does not combine the best features of different witnesses into a single stronger counterexample.

II.217 Genuine initial stages with all six numerical budgets

Source 52, delivered Section 7.

We now remove the upper-length limitation in a different, nonunit-step family. This construction uses the exact spectral cutoff in the source's budget theorem and satisfies all six of its displayed numerical assumptions. Its modulus is composite: the theorem's separate prime hypothesis is explicitly false.

Theorem II.217.1 (A full numerical package over dyadic moduli). *For every sufficiently large dyadic modulus $N = 2^k$, there exist a context of exact density $\alpha = 1/16$, records D, E, F , and a positive integer budget parameter $M = \ell(F.R)$ such that:*

- (i) *Stage 13.4 holds at $\theta = \text{section10Lambda}(\alpha^{32}/16)$, and Stages 13.5 and 13.6 hold;*
- (ii) *the progressions have common step 2 and satisfy $[F.R] \subseteq [E.Q] = [D.P] \subseteq T$, with $\ell(F.R) + 1 \leq \ell(E.Q)$;*
- (iii) *all six numerical assumptions in `section13_bilinear_extraction_of_budgets` hold for M ;*
- (iv) *no weak Stage 13.7 output exists for these records.*

This is not a counterexample to that prime-qualified budget theorem, and is not a claim about its proof-term witnesses. It refutes an all-moduli implication from the specified predicates and numerical conditions to Stage 13.7 for these data.

II.217.1 The six numerical assumptions

Keep c, β, ε as in Section II.212, and write

$$z = \text{section10Zeta}(\alpha^{32}/16), \quad W = 2^{135}\alpha^{-704}, \quad w = 2^{-11q}.$$

The source's natural-number budget parameter is called $\mathfrak{m}$; we call it M to distinguish it from the strip width and $D.m$. Its six hypotheses, read literally from the pinned theorem, are

$$\begin{aligned} 0 < M, \quad M^2 \leq N, \quad M + 1 \leq \ell Q, \\ cN^\beta \leq M, \quad (\ell P)^{-w} \leq z/M, \quad W \leq (cN^\beta)^\varepsilon. \end{aligned} \quad (\text{II.217.1})$$

The theorem also assumes $\alpha \leq 1/6$, genuine initial-stage predicates at the stated cutoff, and primality. We satisfy the density bound and both genuine initial-stage predicates, but not primality. In particular we never describe the construction as satisfying *every hypothesis* of the prime-qualified theorem.

II.217.2 Fixed constants and strict exponent margins

At $\alpha = 2^{-4}$, the Section 10 formulas give

$$\begin{aligned} \theta &= 2^{-37}(2^{-132})^{11/2} = 2^{-763}, & W &= 2^{2951}, \\ K_0 &= 2^{761}, & q &= 2(2K_0 + 1) = 2^{763} + 2, \\ t &= 1882 + 763 \cdot 10477 = 7\,995\,833, & \theta_1 &= 2^{-t}. \end{aligned}$$

In particular $0 < \theta \leq 2^{-130} = \alpha^{32}/4$ and $\alpha \leq 1/6$. Define

$$d = \frac{\theta_1}{64\pi}, \quad c_P = d/2, \quad u = \frac{\theta_1^2}{16q}, \quad v = 2^{-12q}, \quad w = 2^{-11q}, \quad \sigma = uv/2.$$

All these constants are positive and independent of N . One has

$$d < 1, \quad c_P < 1, \quad u < 1/16, \quad 0 < v < w < 1, \quad w \leq 1/2,$$

$$0 < \sigma < u, \quad 2\sigma = uv < 1, \quad \sigma < uw < 1. \quad (\text{II.217.2})$$

The source spectrum bound at 2^{-132} is exactly $2^{74}(2^{-132})^{-10} = 2^{1394} = K$, an integer, so its natural ceiling is K . The corrected denominator in **section10Zeta** is this same K . Consequently

$$z = 2^{-155K}(2^{-132})^{18K}/K = 2^{-2531K-1394} > 0.$$

The other constants retain the values

$$c = 2^{-2532K-1} < 1/2, \quad \beta = 2^{-13 \cdot 2^{9437184}}.$$

The crucial additional margin is $\beta < \sigma$. Since $q < 2^{764}$,

$$\sigma = \frac{2^{-2t-5-12q}}{q} > 2^{-(2t+769+12q)}.$$

Now $2t + 769 = 15\,992\,435 < 2^{24}$ and $12q < 2^{768}$, whence

$$2t + 769 + 12q < 2^{24} + 2^{768} < 2^{769} < 13 \cdot 2^{9437184}.$$

It follows that $\sigma > 2^{-2^{769}} > \beta$. This comparison uses symbolic exponents, not expanded tower values.

II.217.3 Nine sufficient scale conditions

Choose $N = 2^k$ large enough that

$$\begin{aligned} \text{(i)} \quad N &\geq 32, & \text{(ii)} \quad N^u &\geq 2/d, \\ \text{(iii)} \quad N^{1-u} &\geq 16d, & \text{(iv)} \quad N^{u-\sigma} &\geq 2/c_P, \\ \text{(v)} \quad N^\sigma &\geq 2, & \text{(vi)} \quad N &\geq (2K_0)^2, \\ \text{(vii)} \quad N^{uw-\sigma} &\geq c_P^{-w}/z, \\ \text{(viii)} \quad N^{\beta\varepsilon} &\geq Wc^{-\varepsilon}, \\ \text{(ix)} \quad N^{\sigma\varepsilon} &\geq 2^{\varepsilon+1}/\delta. \end{aligned} \quad (\text{II.217.3})$$

Every right side is a fixed positive finite real number. Each growth exponent is strictly positive by (II.217.2), $\beta > 0$, and $\varepsilon > 0$. Therefore one finite integer threshold k_0 makes all nine conditions hold for every $k \geq k_0$. This is a precise existence argument for arbitrarily large finite examples; it does not require searching numerically for their sizes.

Write the strip width as $b = N/16$ in this section, and define

$$p = \lfloor dN^u \rfloor, \quad n = \lfloor N^\sigma \rfloor, \quad M = n.$$

Here p denotes a length, not a prime. Condition (ii) and the floor inequality give

$$c_P N^u \leq p \leq dN^u.$$

Conditions (iii)–(v) then yield

$$2N^\sigma \leq c_P N^u \leq p \leq N/16 = b, \quad N^\sigma/2 \leq n \leq N^\sigma. \quad (\text{II.217.4})$$

In particular $n \geq 1$, $n+1 \leq p$, and $n^2 \leq N^{2\sigma} \leq N$. For $n+1 \leq p$, use $n+1 \leq N^\sigma + 1 \leq 2N^\sigma \leq p$. These inequalities also ensure properness of every progression below.

II.217.4 All initial-stage and Stage 13.6 predicates

Use the common strip context with $T = \{2j : 0 \leq j < b\}$. Put

$$P = Q = (0, 2, p), \quad D.q = q, \quad D.m = p, \quad D.P = P, \quad D.H = [P], \quad E.Q = Q,$$

and use the constant spectral graphs

$$a_{\nu,s} = 0, \quad b_{\nu,s} = \nu N/2 + s, \quad \nu \in \{0, 1\}, \quad |s| \leq K_0.$$

The proof of Lemma II.215.1 works verbatim with the present cutoff: a large coefficient forces $\|2r\|_N \leq 1/(2\theta) = 2^{762} = 2K_0$, and hence belongs to one of these q graphs. Both aliases are retained.

For Stage 13.4, P is proper with nonzero step by $p \leq b$, and $D.H = [P]$. The graph count obeys

$$q < 2^{764} < 2^{t+1526} = 2^{1882+763 \cdot 10479} = Q_{134}(\theta).$$

The exact floor expression is dN^u , so its floor is $D.m = p = \ell P$. Finally all p selected heights are good and

$$\sum_{h \in D.H} C(h) = pC(0) \geq p\alpha^{16}N^{31} \geq \alpha^{32}N^{31}\ell P/8.$$

Together with the uniform spectral cover, this checks every Stage 13.4 conjunct.

For Stage 13.5, $Q = P$, and $p^v/2 \leq p$ since $p \geq 1$ and $v \leq 1$. Every coefficient product with step 2 has centered absolute value at most $2K_0$. Using condition (vi), $p \leq N$, and $w \leq 1/2$, we obtain

$$2K_0 \leq \sqrt{N} \leq N^{1-w} \leq p^{-w}N.$$

This is exactly the required coefficient bound. Every height is strong, hence $\mathcal{H} = [Q]$ has cardinality p and satisfies the critical-height lower bound.

Define

$$F.R = (0, 2, n), \quad F.Y(h) = A \quad (h \in \mathbb{Z}/N\mathbb{Z}).$$

Then $[F.R] \subseteq [Q] \subseteq T$ by (II.217.4). The step and properness requirements hold. The length bound follows from $\beta < \sigma$ and $c < 1/2$:

$$cN^\beta \leq N^\sigma/2 \leq n.$$

The edge-domain identity $F.Y(h) = \mathcal{E}_h = A$ and density $|A| \geq \delta N^2$ are unchanged. The restriction now is

$$Y_h^R = [F.R] \times \mathbb{Z}/N\mathbb{Z}, \quad |Y_h^R| = nN.$$

Its vertical difference is zero and thus ambient affine. The exact aggregate mass is pnN , whereas the two required bounds are δpnN and λpnN . Both hold. This completes Stage 13.6, including the extra upper length and carrier containment stated in the theorem.

II.217.5 All six budgets and the contradiction

The first three conditions of (II.217.1) were proved in (II.217.4); the fourth was just checked. For the fifth, condition (vii) gives

$$np^{-w} \leq N^\sigma (c_P N^u)^{-w} = c_P^{-w} N^{\sigma-uw} \leq z.$$

As $n > 0$, this is precisely $p^{-w} \leq z/n$. For the sixth, condition (viii) implies

$$(cN^\beta)^\varepsilon = c^\varepsilon N^{\beta\varepsilon} \geq W.$$

This verifies the width of the lower target cN^β itself; it is not replaced by the weaker assertion that the chosen n is large. All six numerical assumptions have now been checked with their original constants.

If a weak Stage 13.7 output existed, its first coordinates would lie in $[F.R] \subseteq T$. Lemma II.214.1 gives at most one point in each of the p translated critical rows. The first mass bound therefore implies $L \leq \delta^{-1}$. But condition (ix) and $n \geq N^\sigma/2$ give

$$L \geq n^\varepsilon \geq 2^{-\varepsilon} N^{\sigma\varepsilon} \geq 2/\delta > 1/\delta,$$

a strict contradiction. This proves Theorem II.217.1.

The theorem simultaneously supplies genuine initial stages, the exact cutoff, $R \subseteq Q$, and all six numerical budgets. It does not combine them with the unit input step of Section II.216. That distinction is essential: the strongest properties of these different constructions cannot be silently merged.

[write] **7 October 2026; sources 47–53.** *Checked at the write.* At $\alpha = 1/16$: $43 + 4 \cdot 224 = 939$, $48 + 4 \cdot 256 = 1072$, $100 + 4 \cdot 448 = 1892$, $114 + 4 \cdot 320 = 1394$, $228 + 4 \cdot 576 = 2532$, $2^{20} + 4 \cdot 2^{21} = 9 \cdot 2^{20}$; $t = 1882 + 763 \cdot 10477 = 7,995,833$, $2t + 769 = 15,992,435 < 2^{24}$; in Section II.214, $(k-4)2^{-1892} = 939 + 2^{-1892}$ for $k = 939 \cdot 2^{1892} + 5$; in Section II.215, $t = 1882 + 130 \cdot 10477 = 1,363,892$, $1109 \cdot 2^{1892} - 4 > 0$, and $\lfloor 16/\pi \rfloor = 5$ from $56/11 < 16/\pi < 16/3$. The six budgets are those of `section13_bilinear_extraction_of_budgets` (`Proofs13BilinearExtraction.lean`, line 16, at the pin and at HEAD, where the file has changed elsewhere).

II.218 A prime power family of unit step obstructions

Source 52, delivered Section 8.

The mechanism is not special to powers of two. We record a generalization for the minimal-data construction; the genuine Stage 13.4 certification above remains dyadic.

Theorem II.218.1 (Every fixed prime-power family). *Fix a prime p and an integer s with $p^s \geq 8p$. Put $\alpha = p^{-s}$. For every sufficiently large integer k , with $N = p^k$, there is a context of exact density α and records satisfying Stage 13.5 and 13.6 such that $D.P = E.Q = F.R$ has unit step, but no weak Stage 13.7 output exists.*

Proof. Take $k > s$ and set

$$m = p^{k-s}, \quad T = \{pj \pmod{N} : 0 \leq j < m\}, \quad A = T \times \mathbb{Z}/N\mathbb{Z},$$

$$\phi(x, y) = \lfloor x.\text{val}/p \rfloor \pmod{N}, \quad R = (0, 1, pm).$$

The representatives in T are distinct, $pm \leq N/8 < N$, and $|A| = \alpha N^2$. Horizontal eight-term equality gives $(N/p) \mid \Delta$ for a difference of index sums with

$$|\Delta| \leq 8(m-1) < 8m \leq N/p.$$

Thus $\Delta = 0$. Vertical maps are constant; every arrangement is respected; and the same energy argument gives $C(h) \geq \alpha^{16} N^{31}$ at every height. This verifies the context and strong-height property exactly as before.

Use $D.q = 1$, $D.m = pm$, $D.P = E.Q = F.R = R$, $D.H = [R]$, zero coefficients, and $F.Y(h) = A$. Then $\mathcal{H} = [R]$ and $H = |\mathcal{H}| = pm = \ell Q$. The zero-coefficient Stage 13.5 check applies without change.

For the Stage 13.6 length bound, $K_{13} = 2^{114} \alpha^{-320} \geq 1$ and $0 < \alpha \leq 1$ give

$$0 < c = 2^{-228K_{13}-1} \alpha^{576K_{13}} < \alpha.$$

Also $Q_{13} > 0$ implies $0 < \beta < 1$. Hence $cN^\beta < \alpha N = m \leq pm$. Edge density follows from $\delta = 2^{-43} \alpha^{224} \leq \alpha$. The exact restricted-edge mass is HmN . Its two required bounds reduce to $p\delta \leq 1$ and $p\lambda \leq 1$. These hold since $p\alpha \leq 1/8$ and

$$p\delta = 2^{-43}(p\alpha)\alpha^{223} \leq 2^{-46}, \quad p\lambda = 2^{-48}(p\alpha)\alpha^{255} \leq 2^{-51}.$$

This completes every Stage 13.6 conjunct.

Each point of $[R]$ has a unique form $pj + r$ with $0 \leq r < p$ and $0 \leq j < m$. Its image is j . Two agreements with an ambient affine function in the same residue class r imply

$$(pa - 1)(j - j') = 0 \quad \text{in } \mathbb{Z}/p^k\mathbb{Z}.$$

Since $pa - 1 \equiv -1 \pmod{p}$, it is a unit, so $j = j'$. Thus an affine row has at most p points. A weak output would force

$$L \leq p/\delta, \quad L \geq (pm)^\varepsilon = p^{(k-s+1)\varepsilon}.$$

Here $p, s, \alpha, \delta, \varepsilon$ are fixed and $\varepsilon > 0$. The latter expression grows without bound as k grows. More explicitly, it suffices to choose an integer $k > s$ satisfying

$$(k - s + 1)\varepsilon \log p > \log(p/\delta).$$

Such integers are unbounded, and the inequalities contradict each other. No Stage 13.4 predicate has been used or claimed. $\square$

This family does not assert a counterexample for every composite modulus. It identifies infinitely many counterexamples inside each fixed prime-power family, and locates the exact unit argument needed for that statement.

II.219 A small model and the algebraic boundary

Source 52, delivered Section 9.

The enormous exponents are needed only to force the quantitative Stage 13.7 contradiction. The local-versus-ambient issue is visible on two points.

Proposition II.219.1 (A two-point order-eight model). *In $\mathbb{Z}/18\mathbb{Z}$, let $V = \{0, 2\}$ and $f(0) = 0$, $f(2) = 1$. Then f is an order-eight Freiman homomorphism on V but has no ambient affine extension agreeing at both points.*

Proof. An eight-tuple in V has sum $2r$ when exactly $r \in \{0, \dots, 8\}$ entries are 2. If two such sums agree, $9 \mid (r - s)$. Since $|r - s| \leq 8$, one has $r = s$, and their image sums agree. An affine extension would have intercept zero from its value at 0, while its value at 2 would require $2a = 1 \pmod{18}$, which is impossible. $\square$

At order nine the Freiman property fails: nine copies of 2 and nine copies of 0 have the same domain sum but image sums 9 and 0. The precise order matters. This explanatory example is not a Stage 13.7 counterexample at $N = 18$, and says nothing by itself about an unspecified partition interface.

Lemma II.219.2 (Two-point solvability). *For integer representatives d, e and $N > 0$, the congruence $da = e \pmod{N}$ is soluble if and only if $\gcd(d, N)$ divides e .*

Proof. If $da - e$ is divisible by N , every common divisor of d, N divides e . Conversely, write $g = \gcd(d, N)$ and choose integers u, v with $ud + vN = g$. If $e = gw$, then $a = uw$ solves the congruence. $\square$

Agreement of an affine map at two prescribed points reduces to this lemma after subtraction. For the dyadic strip the output difference is the index difference, whereas the input difference is twice that index difference. Their incompatible 2-adic divisibilities underlie Lemma II.214.1. In the general prime-power construction the same role is played by the extra factor p .

The bounded-order condition cannot force a relation requiring more than eight terms when the domain has been chosen short enough. This is why there is no conflict between exact Freiman preservation and nonextendability. The section map is genuinely linear in its integer parameter, but that parameter need not be given by multiplication by a residue of the ambient ring.

II.220 Source 49's chapter: exact two-cube degeneracy [write]

Added in the ninth write, 7 October 2026.

The next five sections are source 49's Sections 2–6, on the proportion $\delta_{k,1}(p)$ of degenerate ordered pairs of congruent k -dimensional coordinate cubes ($d = 1$ in Gowers's arrangement notation). *Logically they continue source 22's two-cube theorem* (Theorem II.85.1, in Section II.81), in source 22's convention; they are printed here because no number printed earlier may move. Source 49's introduction is Section F.54; its Section 7 is in Part III (Section III.92), its Section 8 in this Part (Section II.226).

What is proved. $\delta_{k,1}(p) = 1 - (1 - p^{-1})^k(1 - |T_k(p)|/p^k)$ exactly, $T_k(p)$ the union of translation loci of nonconstant ternary multiaffine polynomials (Proposition II.221.1); derivative-rank-two polynomials are classified, counted $(62\binom{k}{2} + 84\binom{k}{3} + 24\binom{k}{4})$ and shown to add no locus outside the signed-coordinate arrangement (Theorems II.222.1, II.222.2, Proposition II.222.3); hence $0 \leq \delta_{k,1}(p) - L_k(p) \leq (1 - p^{-1})^k M_k^2/(2p^3)$ and the second coefficient $-k(3k^3 + 14k^2 + 15k - 14)/6$ in every dimension (Theorem II.222.4). $\delta_{k,1} = L_k$ exactly for $k \leq 2$ and, by an exhaustive integer certificate with an algebraic reduction of all large primes, for $k = 3$ at every odd prime (Corollary II.222.5, Theorem II.223.1); the same over $\mathbb{F}_q$ of characteristic at least five (Theorem II.225.2); in characteristic three an exact formula in every dimension from a full-subfield alphabet theorem (Theorem II.225.4, Corollary II.225.5); and for every k , $\delta_{k,1}(p)$ is a fixed Laurent polynomial in p outside a finite set of primes (Theorem II.224.2). This answers source 22's exactness question for $k \leq 3$ (item II.42);

$k \geq 4$ is open. The $d \geq 2$ cross-section statistics of sources 27 and 44 are a different convention and were not reconciled.

Relation to the formal statements (checked at HEAD 06b30ca52). `lemma_15_4` (Sections14_15.lean, line 402) has the exact companion `lemma_15_4_holds` (Proofs15DegenerateArrangements.lean, line 1131), a first-order bound in another convention; nothing in this chapter is formalized.

II.221 Two-cube degeneracy as a translation problem

Source 49, delivered Section 2.

II.221.1 The arrangement convention

Fix an odd prime p and an integer $k \geq 1$. An ordered one-arrangement has parameters

$$h, y^1, y^2 \in \mathbb{F}_p^k, \quad r \in \mathbb{F}_p,$$

and labelled vertices

$$v_\epsilon^j = (y_1^j + \epsilon_1 h_1, \dots, y_k^j + \epsilon_k h_k, r) \quad (j = 1, 2; \epsilon \in \{0, 1\}^k).$$

There are p^{3k+1} arrangements. Coordinates, cubes, and vertices are labelled, and repetitions are allowed. In Gowers' d -arrangement notation this is $d = 1$: the last coordinates satisfy $r_1 = r_2 = r$.

A ternary coefficient pattern is a pair $\eta_1, \eta_2 : \{0, 1\}^k \rightarrow S$, where $S = \{-1, 0, 1\} \subset \mathbb{F}_p$. It is a certificate if

$$\sum_{j=1}^2 \sum_{\epsilon} \eta_j(\epsilon) P(v_\epsilon^j) = 0 \quad \text{for every multiaffine } P \text{ in } k+1 \text{ variables.} \quad (\text{II.221.1})$$

Put $\pi(\epsilon) = (-1)^{|\epsilon|}$. The patterns $(\eta_1, \eta_2) = (c\pi, -c\pi)$, $c \in S$, satisfy (II.221.1) identically in the arrangement parameters. They are the universal patterns in this convention. An arrangement is *degenerate* if it has another certificate. The probability of this event under uniform parameters is $\delta_{k,1}(p)$.

This definition is specific to two cubes. In arrangements with more slices, the vertical moment equations can add independent restrictions; one must not substitute the formulas below for the general d -case.

II.221.2 The exact translation reduction

Let $\mathcal{P}_k(p)$ be the multiaffine polynomials in $\mathbb{F}_p[X_1, \dots, X_k]$ whose values on $\{0, 1\}^k$ lie in S . Multiaffine interpolation on the Boolean cube is invertible over every field, so $|\mathcal{P}_k(p)| = 3^{2^k}$. Define

$$T_k(p) = \bigcup_{\substack{f, g \in \mathcal{P}_k(p) \\ f, g \text{ nonconstant}}} \{t \in \mathbb{F}_p^k : f(X+t) = g(X)\}, \quad (\text{II.221.2})$$

where the equality is a polynomial identity.

Proposition II.221.1 (Exact reduction; prior framework). *For every odd prime,*

$$\delta_{k,1}(p) = 1 - (1 - p^{-1})^k \left(1 - \frac{|T_k(p)|}{p^k} \right). \quad (\text{II.221.3})$$

Proof. If some $h_i = 0$, two labelled vertices in the first cube coincide. Giving them opposite coefficients and all other vertices coefficient zero produces a nonuniversal certificate. Thus every zero-side arrangement is degenerate.

Condition on every $h_i \neq 0$ and put $z_i^j = y_i^j/h_i$ and $t = z^1 - z^2$. The vector t is uniform in $\mathbb{F}_p^k$. Because both last coordinates are r , the moment equations involving the last coordinate are scalar multiples of the horizontal ones. All horizontal squarefree moments vanish exactly when the following polynomial vanishes:

$$\sum_{j,\epsilon} \eta_j(\epsilon) \prod_i (X_i + z_i^j + \epsilon_i).$$

Introduce the slice transform

$$F_\eta(X) = \sum_\epsilon \eta(\epsilon) \prod_i (X_i + \epsilon_i - 1). \quad (\text{II.221.4})$$

At a Boolean vertex v , all summands but $\epsilon = v$ vanish and

$$F_\eta(v) = (-1)^{k-|v|} \eta(v).$$

Thus this transform bijects ternary patterns with $\mathcal{P}_k(p)$ and parity patterns with constant polynomials. Translating X in the horizontal moment polynomial reduces its vanishing to $F_{\eta_1}(X + t) + F_{\eta_2}(X) = 0$. Replacing the second polynomial by its negative gives precisely (II.221.2). If one polynomial is constant the other must be constant as well, and their relation is universal. All other successful identities are nonuniversal. The conditional probability is therefore $|T_k(p)|/p^k$. The probability of nonzero sides is $(1 - p^{-1})^k$, proving the formula. $\square$

The transform differs by a fixed translation from that in [212]; it makes the ternary Boolean values immediate. No approximation has entered the reduction.

II.221.3 Derivative rank and translation cosets

For a multiaffine polynomial put

$$r(f) = \dim \text{span}_{\mathbb{F}_p} \{\partial_1 f, \dots, \partial_k f\}. \quad (\text{II.221.5})$$

The following coset lemma is already proved in [212]. We include a proof because it makes the later codimension counts exact.

Lemma II.221.2 (Translation cosets). *Suppose $p > k$. Then*

$$\text{Stab}(f) = \{u : f(X + u) = f(X)\} = \{u : D_u f = 0\}, \quad D_u = \sum_i u_i \partial_i.$$

This is a linear subspace of dimension $k - r(f)$. For fixed f, g , the translation solution set is empty or one affine coset of this space, and hence has $p^{k-r(f)}$ elements. Successful translation preserves $r(f)$.

Proof. If translation by u fixes f , so does translation by au for every scalar $a \in \mathbb{F}_p$, since the additive group of the prime field is generated by one. The polynomial $f(X + zu) - f(X)$ has z -degree at most $k < p$ and vanishes at every $z \in \mathbb{F}_p$. Each of its coefficients in X is therefore zero. Its coefficient of z is $D_u f$. Conversely, if $D_u f = 0$, the z derivative of $f(X + zu)$ vanishes identically; a polynomial of degree less than p with zero derivative is constant. Rank-nullity proves the dimension assertion. Any two solutions differ by a stabilizer, and translating the partial derivatives preserves their linear dependence relations. $\square$

Lemma II.221.3 (Degree is bounded by derivative rank). *Over any field, a nonconstant multiaffine polynomial satisfies $\deg f \leq r(f)$.*

Proof. Choose a nonzero degree- d monomial $a_A X_A$ with $|A| = d$. In the d derivatives $\partial_j f$, $j \in A$, inspect the d coefficient positions $X_{A \setminus \{i\}}$, $i \in A$. The resulting matrix is diagonal with diagonal entry a_A . For $j \neq i$, the inspected monomial contains X_j , whereas $\partial_j f$ is independent of X_j . The diagonal matrix is nonsingular, proving independence of those derivatives. $\square$

II.221.4 The affine translation arrangement

For $p \geq 5$, an affine member of $\mathcal{P}_k(p)$ is one of three constants; a function of one coordinate with slope ± 1 and either of two intercepts; a function of one coordinate with slope ± 2 and its unique admissible intercept; or a function of two coordinates with slopes in $\{1, -1\}$ and its unique admissible intercept. Consequently the number of rank-zero or rank-one polynomials is

$$n_{\leq 1}(k) = 3 + 6k + 4 \binom{k}{2}. \quad (\text{II.221.6})$$

Here is a useful elementary way to see that the list is exhaustive. For a nonempty proper subset $A \subset \mathbb{F}_p$ and $a \neq 0$, one has $|A \cup (A + a)| \geq |A| + 1$: equality $A + a = A$ would force invariance under all of $\mathbb{F}_p$. Successive nonzero affine slopes therefore give at least one new subset sum until the whole field is reached. A three-valued affine function has at most two nonzero slopes. In the two-slope case, $\{0, a, b, a + b\}$ has at most three elements only if $a = b$ or $a = -b$; fitting the three values into S forces the stated slopes and intercept. The one-slope cases are immediate.

For nonconstant affine f, g the translation equation requires equal slopes. The admissible intercept differences give exactly

$$\mathcal{B}_k(p) = \bigcup_i \{t : t_i \in \{0, 1, -1\}\} \cup \bigcup_{i < j} \{t : t_i = t_j \text{ or } t_i = -t_j\}. \quad (\text{II.221.7})$$

For example, equal single-coordinate slopes 1 give $t_i = 0, \pm 1$; two slopes $1, \pm 1$ give $t_i = \mp t_j$. This affine arrangement and its role in the first-order coefficient are prior results [212]. The single-coordinate witnesses remain valid at $p = 3$, where they already fill the ambient translation space.

If $p > 2k + 1$, successive choice of coordinates gives

$$|\mathbb{F}_p^k \setminus \mathcal{B}_k(p)| = \prod_{i=1}^k (p - 2i - 1). \quad (\text{II.221.8})$$

One excludes $0, \pm 1$ at the first choice and the two signs of each previous coordinate at each later choice. If $p \leq 2k + 1$, there are fewer than k available nonzero sign pairs other than $\{1, -1\}$, so the

complement is empty. The same product is zero, since one of its factors is zero. Thus (II.221.8) in fact holds for every odd prime. Inserting it into (II.221.3) gives the exact lower bound $L_k(p)$ for all odd primes. The question is whether a nonaffine certificate can enlarge this union.

II.222 Rank-two structure and the second degeneracy coefficient

Source 49, delivered Section 3.

II.222.1 A factorization across disjoint coordinate sets

Theorem II.222.1 (Rank-two classification). *Let $p \geq 5$ and $f \in \mathcal{P}_k(p)$ with $r(f) = 2$. Then f is quadratic and has a representation*

$$f(X) = \lambda(L_I(X) + \alpha)(L_J(X) + \beta) + \gamma, \quad (\text{II.222.1})$$

where $\lambda \neq 0$, I, J are disjoint nonempty sets of coordinates, and L_I, L_J have nonzero coefficients exactly on I, J respectively. Each of I, J has at most two elements. If a support has two elements, its two coefficients are equal or opposite.

Proof. By lemma II.221.3, f has degree at most two. A nonconstant affine polynomial has derivative rank one, so f is genuinely quadratic. Write

$$f(X) = \frac{1}{2}X^T H X + b^T X + c,$$

with H symmetric and with zero diagonal. Its derivative rank is $\text{rank}[H \mid b]$. A nonzero symmetric rank-one matrix is a scalar multiple of vv^T ; it cannot have zero diagonal over a field of odd characteristic. Hence $\text{rank } H = 2$ and $b \in \text{im } H$.

Choose i, j with $H_{ij} = \lambda \neq 0$. Their columns are independent, and every column is in their span. Set $u_l = H_{jl}/\lambda$, $v_l = H_{il}/\lambda$. The coefficient of column i in column l is u_l , and that of column j is v_l . Looking at its diagonal entry gives $2\lambda u_l v_l = 0$. Thus $u_l v_l = 0$ for every l , and

$$H = \lambda(uv^T + vu^T).$$

The vectors u, v have disjoint nonempty supports. Since $b \in \text{span}\{u, v\}$, completing the product gives (II.222.1).

Fix a Boolean value of one factor that is nonzero. The other factor's Boolean value set then injects by an affine map into S , and hence has size at most three. The subset-sum growth argument preceding (II.221.7) shows that its support has at most two elements: three nonzero coefficients would already give four values for $p \geq 5$. For two coefficients $a, b \neq 0$, the set $\{0, a, b, a + b\}$ has at most three elements exactly when $a = b$ or $a = -b$. This proves the last assertion. $\square$

II.222.2 No new translation component at rank two

Theorem II.222.2 (Rank-two containment). *If $p \geq 11$, $f, g \in \mathcal{P}_k(p)$, $r(f) = 2$, and $f(X+t) = g(X)$, then $t \in \mathcal{B}_k(p)$.*

Proof. Use theorem II.222.1. Suppose first that I has two coordinates. The Boolean values of $L_I + \alpha$ form a three-term arithmetic progression $A_0, A_0 + a, A_0 + 2a$, $a \neq 0$. For each fixed Boolean choice y in the other block, the corresponding three values of f form an arithmetic progression inside S . For characteristic at least five, a nonconstant such progression has midpoint zero and step 1 or -1 . Therefore

$$C = \{\lambda a(L_J(y) + \beta) : y \in \{0, 1\}^J\} \subset S.$$

Apply the same observation to $g = f(X+t)$. Its three-level step set is $C+w \subset S$, where $w = \lambda a L_J(t)$.

If $|J| = 2$, then C has three distinct elements, so $C = S$. The relation $S + w = S$ forces $w = 0$, since S is a proper subset of the prime field and nonzero translations generate the whole additive group. The two coefficients in L_J are equal or opposite, so this gives $t_j = \pm t_l$ for its two coordinates.

If $|J| = 1$, then C consists of two elements. A two-point subset of S with ordered gap ± 2 has just one placement in S , so $w = 0$. One with ordered gap ± 1 has two placements, and its displacement divided by the gap is 0, 1, or -1 . Since that quotient is the single coordinate t_j , it gives $t_j \in \{0, 1, -1\}$. The assumption $p \geq 11$ in particular rules out wraparound identifications of these small gaps. The argument is symmetric if J is the two-coordinate block.

It remains that both blocks have one coordinate. Write the four integer Boolean values of f as $(A, B, C, D) \in \{-1, 0, 1\}^4$, in the order 00, 10, 01, 11, and set

$$a = A - B - C + D, \quad b = B - A, \quad c = C - A, \quad \Delta = AD - BC = aA - bc. \quad (\text{II.222.2})$$

Thus $f(x, y) = A + bx + cy + axy$. Here $a \neq 0$, $|a| \leq 4$ and $|\Delta| \leq 2$. Under translation (u, v) the coefficient a and the invariant Δ are preserved, while

$$c' = c + au, \quad b' = b + av. \quad (\text{II.222.3})$$

Since $p \geq 11$, equality of these a and Δ invariants modulo p implies equality of their integer representatives. Common negation of f, g lets us take $a > 0$.

Table II.222.1 is the complete list of (c, b) for each (a, Δ) . It can be checked without search over a field: for each entry recover $A = (\Delta + bc)/a$, $B = A + b$, $C = A + c$, $D = A + b + c + a$, and require all four values to lie in $\{-1, 0, 1\}$. The rows contain 31 patterns, the positive- a half of the 62 nonaffine two-variable patterns. For any two entries in a row, the difference divided by a has a coordinate 0, 1, or -1 , or its two coordinates are equal up to sign. This follows directly from the printed list. For the longer $a = 1$ rows, if neither coordinate has absolute value at most one, their absolute values are equal; for $a = 2, 3$ the only additional unequal-size fractions are absent. By (II.222.3), the translation lies in $\mathcal{B}_k(p)$. $\square$

II.222.3 An exact census of rank two

Proposition II.222.3 (Rank-two census). *For every $p \geq 5$, the number of $f \in \mathcal{P}_k(p)$ with derivative rank two is*

$$n_2(k) = 62 \binom{k}{2} + 84 \binom{k}{3} + 24 \binom{k}{4}. \quad (\text{II.222.4})$$

Consequently the exact number of polynomials with rank at least three is

$$M_k = 3^{2^k} - \left(3 + 6k + 66 \binom{k}{2} + 84 \binom{k}{3} + 24 \binom{k}{4} \right). \quad (\text{II.222.5})$$

In particular, $M_1 = M_2 = 0$ and $M_3 = 6258$.

a	Δ	All slope pairs (c, b)
1	-1	$(0, 0), (1, 0), (0, 1), (-1, -1), (-1, -2), (-2, -1)$
1	0	$(0, -1), (-1, 0), (0, 0), (-1, -1)$
1	1	$(1, -1), (-1, 1), (0, -2), (0, -1), (-2, 0), (-1, 0)$
2	-2	$(0, 0), (-2, -2)$
2	-1	$(-1, -1)$
2	0	$(0, -1), (-1, 0), (-1, -2), (-2, -1)$
2	1	$(-1, -1)$
2	2	$(0, -2), (-2, 0)$
3	-1	$(-1, -1), (-2, -2)$
3	1	$(-1, -2), (-2, -1)$
4	0	$(-2, -2)$

Table II.222.1: The entire positive-quadratic-coefficient square table. It is an exact small table in the proof, independently regenerated by the supplied rank-two verifier.

Proof. By theorem II.222.1, essential support has size two, three, or four. For a fixed two-coordinate support, there are $3^4 = 81$ ternary patterns, of which $3 + 12 + 4 = 19$ are affine. Every remaining pattern has rank two, giving 62.

On a three-coordinate support the nonzero Hessian has a unique singleton/pair bipartition. Choose the singleton z in three ways. Up to sign, the centered coordinate on the pair is $w = x + y - 1$ or $w = x - y$. At $z = 0$ and $z = 1$ the function is one of $-1, 0, 1, w, -w$. Conversely, any ordered pair of such restrictions with different w slopes gives an essential rank-two polynomial. There are $5^2 - (3^2 + 1^2 + 1^2) = 14$ choices. With the two choices of w and the three singletons this gives 84.

On four coordinates both factors have three Boolean levels. Normalize their centered values as $U, V \in S$. Write the polynomial as $\theta_0 + \theta_1 U + \theta_2 V + \theta_3 UV$, where $\theta_3 \neq 0$. For each $V \in S$, the three values as U varies lie in S , so their slopes $\theta_1 + \theta_3 V$ lie in S . This is a nonconstant three-term progression, forcing $\theta_1 = 0$ and $\theta_3 = \pm 1$. At $V = 1$ and $V = -1$ the progressions have nonzero slope and hence zero midpoint, which forces $\theta_0 = \theta_2 = 0$. Thus the polynomial is a signed product of two centered pair forms. There are three partitions into pairs, two unsigned forms on each pair, and two choices of overall sign: $3 \cdot 2 \cdot 2 \cdot 2 = 24$.

No pattern is counted twice: its essential support and Hessian bipartition determine the cases, and the forms are distinct up to the sign convention already factored out. Choosing the support gives (II.222.4). Subtracting (II.221.6) and (II.222.4) from 3^{2^k} gives (II.222.5). $\square$

II.222.4 The finite bound and the second coefficient

Theorem II.222.4 (Second-order two-cube degeneracy). *Let $k \geq 1$ and let $p > \max\{7, 2k + 1\}$ be prime. Then*

$$0 \leq \delta_{k,1}(p) - L_k(p) \leq (1 - p^{-1})^k \frac{M_k^2}{2p^3}. \quad (\text{II.222.6})$$

For every fixed k , as $p \rightarrow \infty$ through primes,

$$\delta_{k,1}(p) = \frac{k^2 + 3k}{p} - \frac{k(3k^3 + 14k^2 + 15k - 14)}{6p^2} + O_k(p^{-3}). \quad (\text{II.222.7})$$

Proof. The affine certificates give $\mathcal{B}_k(p) \subseteq T_k(p)$. By theorem II.222.2, rank-two certificates add

no points outside that union. A successful remaining pair has both ranks at least three. By lemma II.221.2, its locus has at most p^{k-3} points. There are M_k^2 ordered pairs. Common negation has no fixed pair and preserves its translation locus, so at most $M_k^2/2$ loci need be charged. A union bound, followed by proposition II.221.1, proves (II.222.6).

For the expansion, the complement product in $1 - L_k$ has k weights equal to one and the k weights $3, 5, \dots, 2k+1$. Their sum is $A = k^2 + 3k$ and the sum of their squares is $k(4k^2 + 12k + 14)/3$. Its second elementary symmetric function is

$$\frac{1}{2} \left(A^2 - \frac{k(4k^2 + 12k + 14)}{3} \right) = \frac{k(3k^3 + 14k^2 + 15k - 14)}{6}.$$

The error in (II.222.6) is already of cubic order, so it cannot change either displayed coefficient. $\square$

[write] **7 October 2026; sources 47–53.** *Against source 22* (checked at the write). Source 22's Theorem II.85.1 proved $L_k(p) \leq \delta_{k,1}(p) \leq L_k(p) + (1 - p^{-1})^k R_k p^{-2}$ with $R_k = \frac{1}{2}(3^{2k} - (2k^2 + 4k + 3))^2$; the rank-two containment moves the error to p^{-3} , with the explicit $M_k^2/2$, so the coefficient of p^{-2} is the one of L_k . The write expanded L_k exactly in $1/p$ for $k = 1, \dots, 8$: the coefficients of p^{-1} and p^{-2} are $k^2 + 3k$ and $-k(3k^3 + 14k^2 + 15k - 14)/6$, and L_1, L_2, L_3 are the polynomials (II.222.8), (II.222.9) and (II.223.1); also $M_3 = 3^8 - (3 + 18 + 198 + 84) = 6258$. Source 22's Question 2 asked whether $\delta_{k,1}(p) = L_k(p)$ for all large p : answered for $k \leq 2$ here and for $k = 3$ by Theorem II.223.1, at every odd prime; open for $k \geq 4$ (item II.42).

The constant $M_k^2/2$ is intentionally explicit and generally large. It is a bound for a union of possible high-rank loci, not an estimate of the actual number of distinct components. The gain is the additional power of p^{-1} and the exact second coefficient. No useful uniform estimate for large k and modest p is inferred from this constant.

Corollary II.222.5 (Exact dimensions one and two). *For every odd prime,*

$$\delta_{1,1}(p) = 4p^{-1} - 3p^{-2}, \quad (\text{II.222.8})$$

$$\delta_{2,1}(p) = 10p^{-1} - 32p^{-2} + 38p^{-3} - 15p^{-4}. \quad (\text{II.222.9})$$

Proof. For $k = 1$ every nonconstant polynomial is affine and its translations give exactly $\mathcal{B}_1$. For $k = 2$ there is no rank at least three; theorem II.222.4 therefore gives equality for $p \geq 11$. For $p = 3, 5$, $\mathcal{B}_2 = \mathbb{F}_p^2$ and $L_2(p) = 1$. At $p = 7$, the only possible unequal nonzero integer quadratic coefficients congruent modulo seven are $3, -4$ and $4, -3$, and their reversals. The Δ values at $a = \pm 3$ are ± 1 , whereas those at $a = \pm 4$ are zero. Hence such aliased coefficients cannot occur in a successful pair. The square table proof still applies to equal integer coefficients, so $T_2 = \mathcal{B}_2$ at seven. Expanding L_1, L_2 proves the formulas. $\square$

II.223 Exactness in dimension three: a finite proof for all odd primes

Source 49, delivered Section 4.

Theorem II.223.1 (Three-dimensional exact count; computer-assisted). *For every odd prime p , $T_3(p) = \mathcal{B}_3(p)$. Consequently*

$$\delta_{3,1}(p) = \frac{18}{p} - \frac{119}{p^2} + \frac{364}{p^3} - \frac{543}{p^4} + \frac{386}{p^5} - \frac{105}{p^6}. \quad (\text{II.223.1})$$

[write] **7 October 2026; sources 47–53.** This is a computer-assisted proof: the reduction is proved in the text and the finite enumeration is `code/49-rank-dual-verify_dimension_three.py`, rerun on a copy at the write with an identical certificate. The intake checked $|T_3(p)|$ against the product formula by brute force at $p = 3, 5, 7, 11, 13$. The two-cube convention here (ordered labelled k -cubes sharing one last coordinate, $d = 1$) is source 22's; sources 27 and 44 treat the cross-section arrangements of Lemma 15.4 for $d \geq 2$, and their second coefficients have not been reconciled with (II.222.7) (not checked at the write).

This theorem depends on the finite enumeration described below and implemented in `verify_dimension_three.py`. We give the complete reduction so the precise content of the computation is visible. The enumerated objects are the $3^8 = 6561$ ternary vertex patterns, not arbitrarily chosen examples of translations or primes.

II.223.1 Cubic translations and a residual gcd

Use integer multiaffine interpolation of each pattern and write

$$F(x, y, z) = h + e_1x + e_2y + e_3z + bxy + cxz + dyz + axyz. \quad (\text{II.223.2})$$

The coefficients satisfy $|a| \leq 8$, $|b|, |c|, |d| \leq 4$ and $|e_i| \leq 2$. For primes $p \geq 17$, equality of cubic coefficients modulo p implies equality as integers. Negating both polynomials reduces the case $a \neq 0$ to $a > 0$.

For two patterns of the same a , use primes on their coefficient symbols for the target pattern. If translation by $(u/a, v/a, w/a)$ is successful, the quadratic coefficients force

$$u = d' - d, \quad v = c' - c, \quad w = b' - b. \quad (\text{II.223.3})$$

The remaining four equations, with denominators cleared, are precisely

$$\begin{aligned} R_1 &= a(e'_1 - e_1) - (bv + cw + vw), \\ R_2 &= a(e'_2 - e_2) - (bu + dw + uw), \\ R_3 &= a(e'_3 - e_3) - (cu + dv + uv), \\ R_0 &= a^2(h' - h) - \{a(e_1u + e_2v + e_3w) + buv + cuw + dvw + uvw\}. \end{aligned} \quad (\text{II.223.4})$$

Success modulo $p \nmid a$ is equivalent to p dividing every residual.

If a rational translation already belongs to $\mathcal{B}_3$ over $\mathbb{Q}$, the same defining equality survives reduction modulo any prime not dividing its denominator. It may be discarded. Thus we retain only pairs for which u, v, w are nonzero, none has absolute value a , and their three absolute values are distinct. This is exactly the rational-complement test.

Lemma II.223.2 (Cubic certificate). *The eight positive- a groups have sizes*

$$(1016, 784, 504, 266, 112, 36, 8, 1), \quad a = 1, \dots, 8.$$

Among their 1 985 589 ordered within-group pairs, precisely 12 288 pass the rational-complement test. The gcd of the four residuals in (II.223.4) has the distribution

<i>gcd</i>	<i>1</i>	<i>2</i>	<i>3</i>	<i>4</i>
<i>number</i>	9408	2496	96	288

In particular no such pair succeeds modulo a prime $p \geq 17$.

Exact finite verification. For each of the 3^8 vertex assignments, apply the Boolean Möbius transform to obtain its integer coefficients. Group by $a > 0$, iterate all ordered pairs in each group, apply the stated test, and compute the gcd of the four displayed integers. This is the entire algorithm. The accompanying verifier performs these integer operations, checks the displayed group sizes and histogram, and fails on a zero gcd or an unexpected value. The histogram lists every surviving pair. A prime at least seventeen divides neither a nor any displayed gcd, which proves the last assertion. $\square$

II.223.2 Quadratic translations and an augmented linear system

If $a = 0$, primes at least seventeen also force the target cubic coefficient to be zero. Equality of the quadratic coefficients modulo such a prime implies equality as integers. Group patterns by their triple $(b, c, d) \neq (0, 0, 0)$. Set

$$H = \begin{pmatrix} 0 & b & c \\ b & 0 & d \\ c & d & 0 \end{pmatrix}, \quad \ell = (e_1, e_2, e_3)^T. \quad (\text{II.223.5})$$

For two patterns in the group, translation is equivalent to

$$\begin{pmatrix} H \\ (\ell + \ell')^T \end{pmatrix} t = \begin{pmatrix} \ell' - \ell \\ 2(h' - h) \end{pmatrix}. \quad (\text{II.223.6})$$

Indeed $Ht = \ell' - \ell$ is the linear-coefficient equation, and the constant equation becomes $2(h' - h) = 2\ell^T t + t^T Ht = (\ell + \ell')^T t$.

Call the 4-by-3 matrix in (II.223.6) M and its right side s . If M has rank three over $\mathbb{Q}$, choose the first nonsingular three-row minor in increasing lexicographic order. Let D be its determinant and let N_i be the corresponding Cramer numerators. The only possible translation is $t = N/D$. Discard it if it belongs rationally to $\mathcal{B}_3$, and otherwise compute the gcd of the four entries of $MN - Ds$.

Lemma II.223.3 (Quadratic certificate). *There are 8788 full-rank ordered pairs in this enumeration. The absolute determinants of their chosen minors have distribution*

$ D $	1	2	3	4	6	8	12	16
number	3264	3008	576	1416	216	252	24	32

Exactly 480 translations pass the rational-complement test; their residual gcds are 3 in 288 cases and 7 in 192 cases.

Exact finite verification. Enumerate the same integer coefficient list, group by (b, c, d) , form the matrix and right side in (II.223.6), and apply the specified minor rule, Cramer's formula, and gcd test. These operations are all on integers. The verifier checks both histograms and the total number of retained pairs. In particular every selected minor remains nonsingular at primes at least seventeen, and no retained pair can satisfy the full system there. $\square$

Why may pairs of rational rank below three be skipped? Suppose a translation is successful in characteristic $p \geq 17$. Modulo the row space of H , the last row $\ell + \ell'$ in M is congruent to 2ℓ , because $\ell' - \ell = Ht$. Thus

$$\text{rank}_{\mathbb{F}_p} M = \text{rank}_{\mathbb{F}_p} [H \mid \ell] = r(f).$$

Here symmetry of H identifies row and column formulations. If $r(f) \leq 2$, the already proved affine or rank-two containment applies. If $r(f) = 3$, M has rank three modulo p , so it also has rank three over $\mathbb{Q}$; reduction cannot increase the rank of an integer matrix. This pair was therefore included in lemma II.223.3. The quadratic computation, together with the cubic computation, covers every possible successful translation outside $\mathcal{B}_3$ for $p \geq 17$.

II.223.3 The remaining primes and the completed proof

At $p = 3, 5, 7$ the signed-coordinate arrangement fills $\mathbb{F}_p^3$, as already observed in (II.221.8). At eleven and thirteen, coordinate permutations and the Boolean reflections $X_i \mapsto 1 - X_i$ preserve $\mathcal{P}_3$. Under reflection the corresponding translation coordinate changes sign. Therefore every translation outside $\mathcal{B}_3$ has an orbit representative whose three coordinates are distinct and increasing in

$$\{2, 3, \dots, (p-1)/2\}.$$

There are four representatives at eleven and ten at thirteen. For each representative the verifier evaluates all 6561 integer polynomials at all eight vertices $t + \epsilon$, reducing modulo p . Exactly three polynomials remain ternary-valued on every shifted vertex: the three constants. Hence there is no nonconstant translation certificate at these representatives, or anywhere outside $\mathcal{B}_3$.

This completes the proof of theorem II.223.1. The polynomial in (II.223.1) is simply the expansion of

$$1 - (1 - p^{-1})^3(1 - 3/p)(1 - 5/p)(1 - 7/p).$$

Remark II.223.4 (What the certificate proves). The large-prime part is a divisibility argument using fixed integer gcds. The two small-prime enumerations close the only remaining cases. Every arithmetic operation used to decide the mathematical claims is exact. This is a computer-assisted finite proof, with independently checkable reductions; it is not a Lean proof, a floating-point optimization, or an assertion about how likely a missed case is. A human compression of the two histograms into a short classification remains a worthwhile project.

II.224 Eventual polynomiality in every fixed dimension

Source 49, delivered Section 5.

The exactness question $T_k = \mathcal{B}_k$ can remain open while the dependence on the prime still has a rigid algebraic form. The classical finite-field method counts complements of rational subspace arrangements by characteristic polynomials; see Athanasiadis [98] and Björner–Ekedahl [216]. Here the translation loci first have to be shown to form one rational affine arrangement with good reduction outside finitely many primes.

Lemma II.224.1 (Rational translation algorithm). *For two multiaffine polynomials $f, g \in \mathbb{Q}[X_1, \dots, X_k]$, the translation locus $f(X + t) = g(X)$ is either empty or a rational affine subspace, computable by finitely many rational linear systems. Moreover, for a fixed finite list of integer polynomials, the same description reduces correctly modulo every prime outside an effectively computable finite set.*

Proof. The empty and constant cases are immediate. First compare the top homogeneous parts. Translation does not change them; if they differ, there is no solution. Suppose they agree in degree

d. Start with $t_0 = 0$ and allowed variations $V = \mathbb{Q}^k$. We describe a descending algorithm. At a stage indexed by j , the current translated f agrees with g in degrees at least j , and every $v \in V$ annihilates, under D_v , its homogeneous parts of degree greater than j . Initially these assertions hold at $j = d$.

For $v \in V$, translation of those higher parts changes nothing: $D_v F = 0$ implies $F(X + v) = F(X)$ in characteristic zero. The change in degree $j - 1$ is therefore exactly $D_v f_j$, where f_j is the current homogeneous part of degree j . Solve the linear system

$$D_v f_j = g_{j-1} - f_{j-1}, \quad v \in V. \quad (\text{II.224.1})$$

If it is inconsistent the translation locus is empty. Otherwise choose one solution v_0 , translate the current polynomial by v_0 , add v_0 to t_0 , and replace V by $V \cap \ker(v \mapsto D_v f_j)$. The required invariant holds for the next stage $j - 1$. Descending through $j = 1$ also matches the constant term. All and only the translations in $t_0 + V$ then solve the identity. This proves the rational description and its algorithm.

Run rational Gaussian elimination with fixed pivot choices at every stage. Record all nonzero pivots, denominator factors, nonzero differences in the initial homogeneous or constant comparisons, and nonzero obstruction entries for inconsistent systems. After excluding their prime divisors, the same operations, ranks, choices of particular solutions, and consistency outcomes hold modulo p . Exclude also $p \leq k$ so that derivative-zero invariance remains valid as in lemma II.221.2. There are finitely many systems for a fixed list of polynomials. Their union of exceptional primes is effective and finite. $\square$

Theorem II.224.2 (Eventual Laurent polynomial). *For every fixed k there are an effectively computable finite set of primes $\mathcal{E}_k$ and a monic polynomial $\chi_k(z) \in \mathbb{Z}[z]$ of degree k such that, for $p \notin \mathcal{E}_k$,*

$$|\mathbb{F}_p^k \setminus T_k(p)| = \chi_k(p), \quad \delta_{k,1}(p) = 1 - \frac{(p-1)^k \chi_k(p)}{p^{2k}}. \quad (\text{II.224.2})$$

Thus the entire large-prime asymptotic expansion is a finite Laurent polynomial, with no terms beyond p^{-2k} . Its coefficients of p^{-1} and p^{-2} are given by (II.222.7).

Proof. Interpolate all ternary Boolean patterns over $\mathbb{Z}$; these give one fixed finite list. Apply lemma II.224.1 to all ordered nonconstant pairs. This gives finitely many rational affine translation loci, each proper since its derivative rank is positive. Compute all their intersections by rational elimination as well, enlarging the exceptional prime set to preserve emptiness and dimension of every intersection after reduction. Outside that set the full intersection pattern is unchanged.

Include the prime two among the exceptional primes. Apply ordinary finite inclusion-exclusion to the complement. Each nonempty intersection of dimension d contributes p^d with an integer sign. The ambient term is p^k ; every other intersection has dimension at most $k - 1$. Hence the result is a single monic integer polynomial of degree k . The second identity follows from proposition II.221.1. The numerator $(p-1)^k \chi_k(p)$ of the subtracted fraction is monic of degree $2k$, so the constant term cancels and only $p^{-1}, \dots, p^{-2k}$ can remain. The coefficient identification follows from theorem II.222.4. $\square$

For $k = 1, 2, 3$, the stronger exact results show that one can take

$$\chi_k(z) = \prod_{i=1}^k (z - 2i - 1)$$

and no odd exceptional primes. In larger dimensions theorem II.224.2 does not prove this factorization. Extra codimension-three or higher loci could change later coefficients while leaving the first two unchanged. Nor is the generic algorithm efficient: the number of ternary patterns is already doubly exponential in k , and naive intersection enumeration adds another major cost. The theorem is a structural finiteness statement and a framework for more selective computations.

II.225 Finite-field extensions and an exact subfield-alphabet theorem

Source 49, delivered Section 6.

The translation formulation also clarifies which parts of the prime-field answer are stable under extension. Define $\mathcal{P}_k(q)$, $T_k(q)$, $\mathcal{B}_k(q)$, and $\delta_{k,1}(q)$ by the same arrangement and ternary certificate definitions over $\mathbb{F}_q$. The exact reduction (II.221.3) remains valid with q in place of p .

II.225.1 Characteristic at least five

Lemma II.225.1 (Rank-two containment over extensions). *Let $q = p^m$ with $p \geq 5$. If $f, g \in \mathcal{P}_k(q)$ are nonconstant, $r(f) \leq 2$, and $f(X + t) = g(X)$, then $t \in \mathcal{B}_k(q)$.*

Proof. Every polynomial in $\mathcal{P}_k(q)$ has coefficients in $\mathbb{F}_p$, because its Boolean values lie in the prime subfield and interpolation uses only addition and subtraction. Its derivative rank is unchanged by extending the scalar field. The affine classification and the rank-two factorization therefore remain valid. In the subset-sum argument, one only needs that a nonzero translation has orbits of size $p > 3$: a nonempty set of size at most three cannot be invariant under it.

For a rank-two factor with a two-coordinate support, the proof of theorem II.222.2 works in every characteristic at least five. A three-term progression with value set $\{-1, 0, 1\}$ has midpoint zero, and the possible placements of a two-element subset are the same as in that proof. Thus it yields a coordinate in $\{0, 1, -1\}$ or two equal coordinates up to sign.

If both factors have singleton supports, translation in those two coordinates is uniquely determined by the two slope equations with nonzero quadratic coefficient. All coefficients are in $\mathbb{F}_p$, so both coordinates of that solution are in $\mathbb{F}_p$. For $p = 5$, every point of $\mathbb{F}_p^2$ belongs to $\mathcal{B}_2(p)$. For $p = 7$ the invariant argument in corollary II.222.5 proves containment. For $p \geq 11$ the square table proves it. The same coordinate equality holds over $\mathbb{F}_q$, and unused coordinates are irrelevant. This completes all cases. $\square$

Theorem II.225.2 (Exact low-dimensional formulas over finite fields). *Let $q = p^m$ with $p \geq 5$. For $k = 1, 2, 3$,*

$$T_k(q) = \mathcal{B}_k(q), \quad \delta_{k,1}(q) = 1 - (1 - q^{-1})^k \prod_{i=1}^k \left(1 - \frac{2i+1}{q}\right). \quad (\text{II.225.1})$$

The expansions in (II.222.8), (II.222.9), and (II.223.1) thus hold with q in place of p .

Proof. Ranks at most two are handled by lemma II.225.1. The only remaining case has $k = 3$ and rank three. Its degree is at most three, which is less than p . The proof of lemma II.221.2 works over $\mathbb{F}_q$ with characteristic $p > \deg f$: invariance under all $a \in \mathbb{F}_p$ gives more roots than the degree in

the auxiliary variable. Rank three gives a trivial stabilizer, so a translation between fixed f, g is unique over $\mathbb{F}_q$.

Since their coefficients lie in $\mathbb{F}_p$, Frobenius sends a solution t to the solution t^p coordinatewise. Uniqueness forces $t^p = t$, and hence $t \in \mathbb{F}_p^3$. The exact prime-field theorem theorem II.223.1 gives $t \in \mathcal{B}_3(p) \subseteq \mathcal{B}_3(q)$. The reverse inclusion is supplied by the affine witnesses. Count the complement by successive choices of sign pairs, as in (II.221.8), and apply the arrangement reduction. $\square$

The all-dimension finite remainder likewise extends: if $p > k$, $p \geq 5$, and $q > 2k + 1$, then

$$0 \leq \delta_{k,1}(q) - L_k(q) \leq (1 - q^{-1})^k \frac{M_k^2}{2q^3}. \quad (\text{II.225.2})$$

Here the characteristic condition, not merely $q > k$, guarantees the translation-coset argument for every degree up to k . The rank-two census is unchanged because scalar extension does not change ranks. Similarly, outside the finite exceptional set in theorem II.224.2, the same rational affine arrangement gives $|\mathbb{F}_{p^m}^k \setminus T_k(p^m)| = \chi_k(p^m)$ for every $m \geq 1$.

Proposition II.225.3 (Polynomiality at a fixed characteristic). *For every odd prime $p > k$, there is a monic integer polynomial $\chi_{k,p}(z)$ of degree k such that*

$$|\mathbb{F}_{p^m}^k \setminus T_k(p^m)| = \chi_{k,p}(p^m) \quad (m \geq 1).$$

This includes primes exceptional for the rational arrangement; the polynomial may then depend on p .

Proof. For fixed f, g with coefficients in $\mathbb{F}_p$, set $K = \ker(u \mapsto D_u f)$ and choose an $\mathbb{F}_p$ -linear complement W . If t is a translation solution over an extension, decompose $t = v + w$ with $v \in K(\mathbb{F}_{p^m})$ and $w \in W(\mathbb{F}_{p^m})$. The stabilizer identity gives $f(X + w) = g(X)$. Frobenius produces the second solution w^p , so $w^p - w \in K(\mathbb{F}_{p^m})$. It is also in $W(\mathbb{F}_{p^m})$, and hence is zero. Thus $w \in \mathbb{F}_p^k$. Every nonempty translation locus is consequently an affine subspace defined over $\mathbb{F}_p$, with an $\mathbb{F}_p$ point, and its extension-field points are exactly the same coset after scalar extension.

There are finitely many pairs. All intersections are affine linear systems over $\mathbb{F}_p$, whose consistency and rank are unchanged by extending the field. Inclusion-exclusion counts their complement by one integer polynomial in p^m , monic of degree k since every locus is proper. $\square$

II.225.2 A full-subfield alphabet

Characteristic three is structurally different because the ternary alphabet is the entire prime subfield. In fact this case has a complete answer in every dimension, as a special case of the following general statement. For a finite subfield $\mathbb{F}_s \subseteq \mathbb{F}_q$, let $\mathcal{P}_{k,s}(q)$ consist of all multiaffine polynomials with Boolean values in $\mathbb{F}_s$, and let $T_{k,s}(q)$ be their nonconstant translation union. The coefficient alphabet is now the full subfield.

Theorem II.225.4 (Exact full-subfield translation arrangement). *Let $q = s^m$ with s a prime power and $m \geq 1$. Then*

$$T_{k,s}(q) = \{t \in \mathbb{F}_q^k : a \cdot t \in \mathbb{F}_s \text{ for some } 0 \neq a \in \mathbb{F}_s^k\}, \quad (\text{II.225.3})$$

and

$$|\mathbb{F}_q^k \setminus T_{k,s}(q)| = \prod_{i=1}^k (q - s^i). \quad (\text{II.225.4})$$

If $k \geq m$, the complement is empty; the product also vanishes. No restriction on the degree relative to the characteristic is needed.

Proof. Boolean interpolation identifies $\mathcal{P}_{k,s}(q)$ with exactly the multiaffine polynomials whose coefficients lie in $\mathbb{F}_s$. If $a \cdot t = c \in \mathbb{F}_s$ with $a \neq 0$, the two affine polynomials $a \cdot X$ and $a \cdot X + c$ certify membership in $T_{k,s}(q)$.

Conversely, suppose a nonconstant f translates to g , both with coefficients in $\mathbb{F}_s$. Let $d \geq 1$ be their top degree, and write f_d for the top homogeneous part. Translation does not change f_d , and its contribution to degree $d-1$ is exactly $D_t f_d$. This identity uses the direct expansion of squarefree monomials and requires no division by factorials. Hence

$$D_t f_d = g_{d-1} - f_{d-1}$$

has all coefficients in $\mathbb{F}_s$. Choose a top monomial $c_A X_A$ with $c_A \neq 0$ and $i \in A$. The coefficient of $X_{A \setminus \{i\}}$ in $D_t f_d$ is an $\mathbb{F}_s$ -linear form $a \cdot t$ whose t_i coefficient is $c_A \neq 0$. It lies in $\mathbb{F}_s$, giving (II.225.3).

Let $V = \mathbb{F}_q/\mathbb{F}_s$, an $(m-1)$ -dimensional vector space over $\mathbb{F}_s$. A point is outside the union exactly when the classes of $t_1, \dots, t_k$ are linearly independent in V . Their ordered choices number

$$\prod_{j=0}^{k-1} (q/s - s^j),$$

when $k < m$, and every class has s lifts in $\mathbb{F}_q$. Multiplying by s^k proves (II.225.4). If $k \geq m$, independence is impossible and the product contains the zero factor $q - s^m$. $\square$

Corollary II.225.5 (All-dimensional ternary count in characteristic three). *For every $q = 3^m$ and every $k \geq 1$,*

$$\delta_{k,1}(q) = 1 - (1 - q^{-1})^k \prod_{i=1}^k \left(1 - \frac{3^i}{q}\right). \quad (\text{II.225.5})$$

In particular, every arrangement is degenerate when $k \geq m$.

Proof. In characteristic three, $\{-1, 0, 1\} = \mathbb{F}_3$. Apply theorem II.225.4 with $s = 3$ and the exact arrangement reduction. $\square$

The distinction from (II.225.1) is essential. For instance, in $\mathbb{F}_9$ choose $a \notin \mathbb{F}_3$. The polynomials $X + Y$ and $X + Y + 1$ have ternary Boolean values, and translation by $(a, 1 - a)$ sends the first to the second. Neither coordinate is in $\{0, 1, -1\}$, and the two coordinates are neither equal nor opposite. Thus this translation is outside $\mathcal{B}_2(9)$. In fact corollary II.225.5 says that every two-dimensional arrangement over $\mathbb{F}_9$ is degenerate. This example explains why the characteristic restriction on the signed-coordinate formula cannot simply be removed.

II.226 Sources 50 and 49: Lemma 16.10 is false as catalogued, and an acyclic replacement [write]

Added in the ninth write, 7 October 2026.

The next nine sections are source 50's Sections 2–9 and source 49's Section 8. *Logically they continue the Section 16 material* of sources 19 (sampling and the audit of the closing comparison,

Section II.65), 23 (gluing, Section II.81) and 39 (shared anchors, Section II.151); they are printed here because no number printed earlier may move. Source 50’s introduction is Section F.55.

What source 50 proves. For a word a with symbol masses $b_{(1)} \geq \dots \geq b_{(R)}$ on a progression, q polynomials of degree at most d cover at most $\max_s (B_s + (q - s)dR)$ points, exactly B_q when $b_{(q)} \geq dR$, with rigidity and a deficit for every nonconstant list (Theorem II.228.2); the law lifts to product boxes and adaptive partitions (Theorem II.229.1, Corollary II.229.2); words balanced on every long progression exist and are constructed deterministically (Theorems II.230.2, II.230.3), and every branch profile is realized (Theorem II.231.1). *Application:* at $k = \theta = \gamma = 1$ and every prime $p > 2^{4096Q^4}$, $Q = 2^{2^{1024}}$, the full-domain function $F(h, y) = a(y)$ of an $R = 4Q$ -valued balanced word satisfies both packaged premises of the catalogue’s `lemma_16_10`, while every cover allowed by its conclusion retains at most $5p^2/16$ of the $p^2/2$ points required (Theorem II.233.1). The catalogue statement is therefore false; this is not a counterexample to Gowers’s Lemma 16.10 in context, whose full-domain product property at $\gamma = 1$ is equivalent to multiaffinity and excludes these witnesses (Proposition II.234.1). It strengthens source 19’s finding that the displayed closing comparison fails.

What source 49’s Section 8 proves. From seven stated interfaces I1–I7 (earlier statements of the paper in a corrected form), an acyclic cover theorem with four separate controls N_j , Q_j , $c_{j,D}$, $T_{j,D}$, the global core chosen before the local tolerance (Theorem II.235.5), and a recursive replacement for the conclusion of Corollary 16.11 (Corollary II.235.11), with width $p^{\exp(-(C/\alpha)^C)}$ above $\exp(\exp((C/\alpha)^C))$. It reproduces source 19’s audit and does not certify I1–I7. Source 50’s counterexample is why such a replacement must change the controls.

Relation to the formal statements (checked at HEAD 06b30ca52). `lemma_16_10` (`Section16.lean`, line 851) is open in the ledger and *formally refuted*: `lemma_16_10_packaged_premises_counterexample` (`Proofs16PackagedPremisesCounterexample.lean`, line 13), along source 50’s route; the endpoint characterization is `section16_unit_product_iff_multilinear` (`Proofs16MultilinearProduct.lean`, line 89); `corollary_16_11` (`Section16.lean`, line 893) is open, proved conditionally on the Section 16 induction. Sources 49 and 50 are named in `FORMALIZATION_STATUS.txt` as the two candidates for a repaired statement. Nothing else in this chapter is formalized.

II.227 Progressions, polynomial lists, and branch masses

Source 50, delivered Section 2.

Definition II.227.1 (Proper progression and box). For a prime p , a nonempty proper modular arithmetic progression is

$$I = \{a + bt : 0 \leq t < n\} \subseteq \mathbb{F}_p, \quad b \neq 0, \quad 1 \leq n \leq p.$$

Its length is $n = |I|$. A product box is $I_1 \times \dots \times I_D$, and its width is $\min_i |I_i|$. A common-step box has the same difference in all axes. The repository uses common-step boxes. Our estimates allow independent steps and hence apply to that narrower class.

Singletons can also be represented with difference zero. This harmless representational redundancy never affects the long-axis estimates. The empty box is allowed when discussing universal cover predicates; its width is zero, and it contributes no mass. All counting arguments concern nonempty cells. Properness is essential: a formal length with repeated residues cannot be used in place of the number of distinct points.

Definition II.227.2 (Cover by a polynomial list). Let $I \subseteq \mathbb{F}_p$ be nonempty, let $a : I \rightarrow S \subseteq \mathbb{F}_p$, and let d, q be nonnegative integers. Define the unnormalized optimum

$$\text{Cov}_{q,d}(a; I) = \max_{P_1, \dots, P_q \in \mathbb{F}_p[Y], \deg P_i \leq d} |\{y \in I : a(y) \in \{P_1(y), \dots, P_q(y)\}\}|.$$

For $q = 0$ the optimum is zero. Repetitions are permitted. A list of fewer than q polynomials can be padded without reducing its covered set.

Write $S = \{v_1, \dots, v_R\}$ with distinct v_j . The symbol counts are

$$b_j(I) = |\{y \in I : a(y) = v_j\}|.$$

When order matters, write $b_{(1)} \geq \dots \geq b_{(R)}$ and $B_s = b_{(1)} + \dots + b_{(s)}$, with $B_0 = 0$. A *useful constant* is a constant polynomial whose value belongs to S . We count distinct useful constants in a list; repeated copies do not create additional branches.

For a probability vector $\pi = (\pi_1, \dots, \pi_R)$, set

$$M_s(\pi) = \sum_{j=1}^s \pi_{(j)}, \quad M_0(\pi) = 0.$$

The sequence M_s records the greatest mass that can be retained by selecting s branches. It is the relevant complexity profile in this article. It is not being identified with Shannon entropy or with the rank of a polynomial.

Definition II.227.3 (Multiaffine polynomials). A polynomial on $\mathbb{F}_p^D$ is multiaffine if its degree in each coordinate is at most one. Equivalently, it has the form

$$\mu(z) = \sum_{J \subseteq \{1, \dots, D\}} c_J \prod_{j \in J} z_j.$$

Constants and lower-dimensional terms are included. This is the meaning of `IsMultilinear` in the repository [219].

In particular, fixing every coordinate except the last gives an affine polynomial in the last coordinate. This elementary observation is the only property of multiaffinity needed for the main counterexample.

II.228 Exact one-dimensional cover laws

Source 50, delivered Section 3.

Lemma II.228.1 (Finite-alphabet root bound). *If $P \in \mathbb{F}_p[Y]$ is nonconstant of degree at most d and $a : I \rightarrow S$ with $|S| = R$, then*

$$|\{y \in I : P(y) = a(y)\}| \leq dR.$$

Proof. The agreement set is contained in the union, over $v \in S$, of the roots of $P - v$. Each $P - v$ is a nonzero polynomial of degree at most d and therefore has at most d roots. Summing these R bounds proves the claim. The usual root bound follows by factoring out $Y - y_0$ at each distinct root, so no restriction on the relative sizes of d and p is needed for the inequality. The bound is simply uninformative when $dR \geq |I|$. $\square$

Theorem II.228.2 (Finite cover envelope and exact optimum). *Let $a : I \rightarrow S$, $|I| = n$, $|S| = R$, and $d \geq 1$. Then*

$$B_{\min(q,R)} \leq \text{Cov}_{q,d}(a; I) \leq \min \left\{ n, \max_{0 \leq s \leq \min(q,R)} (B_s + (q-s)dR) \right\}. \quad (\text{II.228.1})$$

If $1 \leq q \leq R$ and $b_{(q)} \geq dR$, then

$$\text{Cov}_{q,d}(a; I) = B_q. \quad (\text{II.228.2})$$

If $b_{(q)} > dR$, every optimal list of size q consists of q distinct useful constants whose total symbol mass is B_q .

Proof. The $\min(q, R)$ heaviest useful constants establish the lower bound. Consider any list and let s be its number of distinct useful constants. They cover at most B_s points. Each remaining nonconstant polynomial covers at most dR points by Lemma II.228.1; an outside-alphabet constant or an unneeded duplicate covers no additional points and can be charged the same upper bound. The union has size at most $B_s + (q-s)dR$. This proves (II.228.1).

When $q \leq R$, for every $s \leq q$,

$$B_q - B_s = \sum_{j=s+1}^q b_{(j)} \geq (q-s)b_{(q)}.$$

Consequently $B_s + (q-s)dR \leq B_q$ under the stated hypothesis. The lower and upper bounds agree. If $b_{(q)} > dR$ and $s < q$, the inequality is strict. An optimal list must therefore have $s = q$; all its entries are distinct useful constants, and their masses must sum to B_q . $\square$

Corollary II.228.3 (Quantitative optimizer deficit). *Under $1 \leq q \leq R$ and $b_{(q)} > dR$, a list with s distinct useful constants misses at least*

$$(q-s)(b_{(q)} - dR) \quad (\text{II.228.3})$$

points relative to the optimum B_q . Thus a list covering at least $B_q - D$ points has

$$q-s \leq \frac{D}{b_{(q)} - dR}.$$

Proof. Subtract its upper bound in (II.228.1) from B_q and use the preceding proof. $\square$

Remark II.228.4 (What is and is not sharp). The exact-optimum conclusion, and hence the leading branch-mass profile, is sharp because the heaviest constant graphs attain it. We do not claim that the root envelope in (II.228.1) is attained for every list, every word, or every n, d, R . Its utility is that the envelope becomes an exact answer once the individual branch masses exceed the root-counting allowance.

Corollary II.228.5 (Prescribed approximate masses). *Suppose $|b_j/n - \pi_j| \leq \varepsilon$ for every symbol. For $q \leq R$,*

$$M_q(\pi) - q\varepsilon \leq \frac{\text{Cov}_{q,d}(a; I)}{n}, \quad (\text{II.228.4})$$

$$\frac{\text{Cov}_{q,d}(a; I)}{n} \leq \min \left\{ 1, \max_{0 \leq s \leq q} \left(M_s(\pi) + s\varepsilon + (q-s)\frac{dR}{n} \right) \right\}. \quad (\text{II.228.5})$$

For the uniform profile and $n \geq dR^2$ this gives

$$\frac{\text{Cov}_{q,d}(a; I)}{n} \leq q \left(\frac{1}{R} + \varepsilon \right). \quad (\text{II.228.6})$$

If $n(1/R - \varepsilon) \geq dR$, the exact value is B_q/n .

Proof. For any set of s symbols its mass is at most $M_s(\pi) + s\varepsilon$. Selecting the q symbols of largest prescribed mass gives the lower bound. Apply Theorem II.228.2. In the uniform case both $1/R + \varepsilon$ and dR/n are bounded by $1/R + \varepsilon$, giving (II.228.6). The last assertion follows from $b_{(q)} \geq n(1/R - \varepsilon)$. $\square$

II.229 Product boxes, adaptive partitions, and global rigidity

Source 50, delivered Section 4.

II.229.1 A dimension-free lifting of the obstruction

The following statement deliberately allows far more than multiaffine graphs. The coefficients of the final-coordinate polynomial may depend arbitrarily on the other coordinates.

Theorem II.229.1 (Product-box cover law). *Let T be any finite nonempty set, $I \subseteq \mathbb{F}_p$, and $F(h, y) = a(y)$ on $T \times I$. Suppose each candidate graph $\mu_i(h, y)$ is, for every fixed h , a polynomial of degree at most d in y . Then the maximum number of covered points using q graphs is exactly*

$$|T| \text{Cov}_{q,d}(a; I). \quad (\text{II.229.1})$$

In particular, every list covers at most $|T|$ times the upper envelope in (II.228.1). If $1 \leq q \leq R$ and $b_{(q)} \geq dR$, this maximum is $|T|B_q$, and the maximum remains the same when the graphs are required to be multiaffine in all coordinates of a product box $T \times I$.

Proof. Fix h and apply the definition of $\text{Cov}_{q,d}$. Sum over $h \in T$ for the upper bound. A maximizing univariate list, used independently of h , attains (II.229.1) in the stated unrestricted class. Under $b_{(q)} \geq dR$, the heaviest constants attain it; those constants are multiaffine as well. $\square$

Corollary II.229.2 (Adaptive partition bound). *Suppose $a : \mathbb{F}_p \rightarrow S$ is (π, ε) -balanced on every proper modular progression of length at least L . Let a nonempty finite family of pairwise disjoint nonempty product boxes $C_u = T_u \times I_u$ have $|I_u| \geq L$. On each cell choose at most q graphs with final-coordinate degree at most d , allowing the partition and the lists to depend on the entire word. The covered fraction of the union is at most*

$$\min \left\{ 1, \max_{0 \leq s \leq \min(q, R)} \left(M_s(\pi) + s\varepsilon + (q - s) \frac{dR}{L} \right) \right\}. \quad (\text{II.229.2})$$

For a uniform profile, $q \leq R$, and $L \geq dR^2$, it is at most $q(1/R + \varepsilon)$.

Proof. Apply Corollary II.228.5 and Theorem II.229.1 separately to each cell, and sum their cardinality bounds. Disjointness turns this sum into the total covered mass. All estimates were established simultaneously for every allowable final axis, so no independence between the word, the partition, and the lists is being assumed. $\square$

No bound on the number of cells is needed. The other axes can be arbitrarily short, the common differences can change between cells, and the lists may change between fibres. Requiring every axis to be long or the differences to agree only restricts the adversary further.

II.229.2 Global polynomial stability

The preceding fibrewise theorem is sufficient for the counterexample. When the candidate graphs are polynomial in the base coordinates as well, one can also control how many *globally* nonconstant graphs occur in a nearly optimal list.

Lemma II.229.3 (Rectangular zero bound). *Let $P \in \mathbb{F}_p[X_1, \dots, X_D]$ be nonzero and satisfy $\deg_{X_i} P \leq e_i$. For nonempty finite sets $A_i \subseteq \mathbb{F}_p$,*

$$\frac{|\{z \in \prod_i A_i : P(z) = 0\}|}{\prod_i |A_i|} \leq \sum_{i=1}^D \frac{e_i}{|A_i|}. \quad (\text{II.229.3})$$

Proof. Induct on D . In dimension one this is the root bound. Write P as a polynomial in X_D and choose its nonzero leading coefficient P_0 . Outside the zero set of P_0 on the base grid, the last-coordinate polynomial has at most e_D roots. On the exceptional base points allow all last coordinates. The induction hypothesis bounds the exceptional fraction by $\sum_{i < D} e_i/|A_i|$. Adding $e_D/|A_D|$ proves the result. $\square$

Theorem II.229.4 (Global-graph envelope and rigidity). *Let $C = A_1 \times \dots \times A_k \times I = T \times I$, and put $F(h, y) = a(y)$. Consider polynomial graphs with base-coordinate degrees at most e_i and final-coordinate degree at most $d \geq 1$. Write $\lambda_j = b_{(j)}/|I|$ and*

$$z_C = \sum_{i=1}^k \frac{e_i}{|A_i|} + \frac{dR}{|I|}. \quad (\text{II.229.4})$$

Any list of $q \leq R$ graphs with s distinct globally useful constants covers at most the following fraction of C :

$$\sum_{j=1}^s \lambda_j + (q - s)z_C. \quad (\text{II.229.5})$$

If $\lambda_q \geq z_C$, the optimum is $\sum_{j=1}^q \lambda_j$. If $\lambda_q < z_C$, a list whose covered fraction is within Δ of the optimum satisfies

$$q - s \leq \frac{\Delta}{\lambda_q - z_C}. \quad (\text{II.229.6})$$

In particular, every optimal list consists of globally constant branches.

Proof. For a globally nonconstant polynomial μ and each occupied symbol v_j , the polynomial $\mu - v_j$ is nonzero. Apply Lemma II.229.3 on $T \times I_j$, where $I_j = \{y \in I : a(y) = v_j\}$. It gives at most

$$|T| \left(|I_j| \sum_i \frac{e_i}{|A_i|} + d \right)$$

agreements on that subgrid. Summing over occupied symbols gives at most $|C|z_C$ agreements. A constant outside the alphabet contributes zero. The s distinct useful constants contribute at most $|C|\sum_{j \leq s} \lambda_j$, so the union bound proves (II.229.5). The rest follows by subtracting this bound from the mass of the q heaviest constants, exactly as in Theorem II.228.2. $\square$

For multiaffine graphs and width at least w , one may use $z_C \leq (k + R)/w$. Thus on sufficiently wide balanced boxes, replacing any constant branch by a genuinely nonconstant global polynomial creates a fixed loss in covered mass. This strengthens an order-of-magnitude obstruction into a classification of exact optimizers and a stability estimate.

Corollary II.229.5 (Stability across a partition). *Under Theorem II.229.4, suppose all cells have $\lambda_{q,u} - z_{C_u} \geq g > 0$. If a cellwise list loses total mass at most $\Delta \sum_u |C_u|$ relative to the sum of the cellwise optima, then*

$$\frac{\sum_u |C_u|(q - s_u)}{\sum_u |C_u|} \leq \frac{\Delta}{g}.$$

Proof. Multiply the per-cell deficit bound by the cell cardinality and sum. □

II.230 Balanced words at logarithmic scales

Source 50, delivered Section 5.

II.230.1 A self-contained concentration estimate

Lemma II.230.1 (Bernoulli concentration). *If $X_1, \dots, X_n$ are independent Bernoulli variables with common mean u , then*

$$\mathbb{P} \left(\left| n^{-1} \sum_i X_i - u \right| > \varepsilon \right) \leq 2 \exp(-2n\varepsilon^2).$$

Proof. For a Bernoulli variable X , let $g(t) = \log \mathbb{E} \exp(t(X - u))$. Then $g(0) = g'(0) = 0$, and $g''(t)$ is the variance of a Bernoulli variable under exponential tilting. Hence $0 \leq g''(t) \leq 1/4$ and $g(t) \leq t^2/8$. Independence gives

$$\mathbb{E} \exp \left(t \sum_i (X_i - u) \right) \leq \exp(nt^2/8).$$

Markov's inequality with $t = 4\varepsilon$ bounds the upper tail by $\exp(-2n\varepsilon^2)$. Apply the same argument to $-\sum_i (X_i - u)$ for the lower tail, and add the bounds. The degenerate cases $u = 0, 1$ are immediate. This is the Bernoulli specialization of Hoeffding's inequality [217]. □

Theorem II.230.2 (Simultaneous progression balance). *Let $p \geq R$ be prime, let $S \subseteq \mathbb{F}_p$ have size R , and let π be a probability vector on S . For $\varepsilon > 0$, put*

$$L(p, R, \varepsilon) = \left\lceil \frac{\log(4Rp^3)}{2\varepsilon^2} \right\rceil. \tag{II.230.1}$$

There exists a word $a : \mathbb{F}_p \rightarrow S$ such that

$$\left| \frac{|\{y \in I : a(y) = v_j\}|}{|I|} - \pi_j \right| \leq \varepsilon \tag{II.230.2}$$

for every symbol and every proper modular progression I of length at least L . If $L > p$, the assertion about such progressions is vacuous; all subsequent applications explicitly ensure $L \leq p$.

Proof. Assign the symbols independently with probabilities π_j . There are at most p^3 triples consisting of a start, a nonzero difference, and a length. For one such progression and one symbol, Lemma II.230.1 bounds the failure probability by $2e^{-2\varepsilon^2|I|}$. The union bound is at most

$$2Rp^3e^{-2\varepsilon^2L} \leq \frac{1}{2}.$$

A word with no failures therefore exists. Repeated parameterizations of the same set only overcount the bad events, which is harmless. $\square$

II.230.2 Deterministic construction with rational probabilities

Theorem II.230.3 (Exact conditional-expectation construction). *If π and ε are rational and the expected number of balance failures is less than one, a balanced word can be constructed deterministically using polynomially many exact arithmetic operations in p and R . A direct implementation uses $O(R^2p^4 + Rp^3)$ arithmetic operations after binomial-tail precomputation, with intermediate integers of bit length polynomial in p , $\log R$, and the input denominator lengths.*

Proof. Enumerate the progression–symbol upper and lower bad events. At most $2Rp^3$ events are needed. After fixing an initial segment of the word, an event’s conditional probability is a binomial tail: the already fixed count is known, and the unfixed positions of its progression remain independent Bernoulli variables. For rational π_j these probabilities are rational and can be computed exactly.

Let Z be the number of bad events. Initially $\mathbb{E}Z < 1$. At a new position, the current conditional expectation is the π -weighted average of the conditional expectations after the possible assignments. Choose an assignment of positive probability that does not increase that expectation. Once every position has been assigned, the conditional expectation is the integer Z , still less than one. Thus $Z = 0$.

For completeness, precompute the binomial cumulative distribution tables for all remaining lengths $0 \leq m \leq p$ and all symbols. With a common probability denominator D , every event probability can be represented on denominator D^p . These tables and powers require polynomially many exact operations. An event can then be evaluated from its fixed count and remaining length by two table lookups. Evaluating all events for all R candidates at each of p positions gives $O(R^2p^4)$ operations. Progression membership is determined by solving $y = a + bt$ and checking $0 \leq t < n$ in the standard representatives, so maintaining the fixed counts also has polynomial cost. The scaled integer scores have $O(p \log D + \log(Rp^3))$ bits. This proves the claimed arithmetic bound; it is a bound in the field size, not in $\log p$. $\square$

The algorithm can equivalently stop at the lexicographically first balanced word found by exhaustive search. Conditional expectation is included to show that this is not merely an unstructured existence argument. It does not make the astronomical instance in Section II.233 practical.

II.230.3 An explicit threshold for power-size cells

Lemma II.230.4 (A convenient threshold). *Let $R \geq 2$, $0 < c \leq 1$, and $\varepsilon = 1/(4R)$. If*

$$\log p \geq \frac{128R^2}{c^2}, \tag{II.230.3}$$

then

$$p^c \geq L(p, R, \varepsilon), \quad p^c \geq R^2, \quad p \geq R.$$

Proof. Set $t = \log p$. The assumption implies $t \geq \log(4R)$ and $t \geq 1$. Hence

$$L = \left\lceil 8R^2(\log(4R) + 3t) \right\rceil \leq 32R^2t + 1 \leq 33R^2t.$$

The power-series bound $e^{ct} \geq (ct)^2/2$, together with $c^2t \geq 128R^2$, gives $p^c \geq 64R^2t \geq L$ and $p^c \geq R^2$. Since $c \leq 1$, $p \geq p^c \geq R^2 \geq R$. $\square$

II.231 Realizing arbitrary local cover profiles

Source 50, delivered Section 6.

Theorem II.231.1 (Asymptotically exact profile realization). *Fix $R, d \geq 1$, a probability vector π on R symbols, and a length function $L_p \leq p$ with $L_p/\log p \rightarrow \infty$ along the primes. There are words $a_p : \mathbb{F}_p \rightarrow S_p$, $|S_p| = R$, such that, for every $q \leq R$,*

$$\sup_{I: |I| \geq L_p} \left| \frac{\text{Cov}_{q,d}(a_p; I)}{|I|} - M_q(\pi) \right| \rightarrow 0. \quad (\text{II.231.1})$$

The same limit is valid for the best covered fraction of any nonempty union of disjoint nonempty product boxes whose final axes have length at least L_p , allowing an adaptive choice of the boxes and at most q final-degree- d graphs on each cell.

Proof. Set

$$\varepsilon_p = \sqrt{\frac{\log(4Rp^3)}{2L_p}}.$$

The proof of Theorem II.230.2, with the integer minimum length $\lceil L_p \rceil$, supplies words balanced within ε_p , and $\varepsilon_p \rightarrow 0$. Corollary II.228.5 implies, uniformly in the progression,

$$M_q(\pi) - q\varepsilon_p \leq \frac{\text{Cov}_{q,d}}{|I|} \leq M_q(\pi) + q\varepsilon_p + \frac{qdR}{L_p}.$$

The upper inequality uses $M_s \leq M_q$ for $s \leq q$. Both errors tend to zero. For boxes, apply Corollary II.229.2; using the q globally heaviest prescribed constants in every cell gives the matching lower bound. $\square$

If $\pi_{(q)} > 0$, eventually $|I|(\pi_{(q)} - \varepsilon_p) \geq dR$ on every admissible axis. The univariate optimum is then *exactly* the mass of the q empirically heaviest symbols, not merely asymptotic to it. With long base axes as well, Theorem II.229.4 classifies the global polynomial optimizers.

Corollary II.231.2 (No branch-budget reset). *Fix $q < R$, a loss $\eta \in (0, 1)$, a degree bound d , and any $c > 0$. If $q/R < 1 - \eta$, then, for all sufficiently large primes, there is a full-domain function of the form $F(h, y) = a(y)$ such that:*

- (i) *every fixed- y cross-section is constant;*
- (ii) *every final-coordinate fibre is exactly covered by R constant graphs;*
- (iii) *no partition into boxes of width at least p^c admits a cellwise cover by q degree-at-most- d final-coordinate graphs retaining a fraction $1 - \eta$ of the full box.*

Proof. Use the uniform profile and Theorem II.231.1 with $L_p = p^c$ when $c \leq 1$. For $c > 1$ no nonempty proper box of that width exists, so the nonexistence conclusion is automatic. Choose the error smaller than the fixed gap $1 - \eta - q/R$. $\square$

Remark II.231.3 (Why a power-width conclusion matters). The condition $L_p/\log p \rightarrow \infty$ is sufficient, not claimed optimal. It includes every fixed positive power p^c , however small the exponent is. Thus replacing a failed width exponent by a smaller positive constant cannot by itself remove a branch-count obstruction. Allowing singleton cells would, of course, change the problem completely.

II.231.1 A stronger scale for growing alphabets

The deterministic root allowance dR is appropriate for arbitrary words, but can be reduced substantially when the word itself is being constructed. The following result gives a second, affirmative quantitative refinement.

Theorem II.231.4 (Simultaneous avoidance of nonconstant polynomial graphs). *Let p be prime, $2 \leq R \leq p$, $1 \leq d < p$, and $0 < \tau < 1$. Put*

$$H = (d + 1) \log p + \log 4, \quad (\text{II.231.2})$$

$$T = \left\lceil \max\{e^2 d, H\} \right\rceil, \quad L_{\text{bal}} = \left\lceil \frac{4R}{\tau^2} \log(8Rp^3) \right\rceil. \quad (\text{II.231.3})$$

For every alphabet $S \subseteq \mathbb{F}_p$ of size R , there exists a word $a : \mathbb{F}_p \rightarrow S$ with the following simultaneous properties:

- (i) *On every proper progression I with $|I| \geq L_{\text{bal}}$, each symbol has between $(1 - \tau)|I|/R$ and $(1 + \tau)|I|/R$ occurrences.*
- (ii) *Every nonconstant polynomial of degree at most d agrees with a on fewer than T points of every proper progression.*

Consequently, on every progression with

$$|I| \geq \max \left\{ L_{\text{bal}}, \frac{R(T + 1)}{1 - \tau} \right\}, \quad (\text{II.231.4})$$

the exact optimum for each $q \leq R$ is the mass of the q heaviest constants, and every optimal list contains only useful constants. In particular, for fixed τ the sufficient scale is $O(R(d + 1) \log p)$, rather than the R^2 -type scale obtained by combining absolute symbol balance with the worst-case root bound. These conclusions are nonvacuous when the displayed scale is at most p .

Proof. Choose the symbols independently and uniformly. For one symbol count on an n -point progression, the mean is $\mu = n/R$. The exponential-moment proof of the Bernoulli tail estimate gives

$$\mathbb{P}(|X - \mu| > \tau\mu) \leq 2e^{-\tau^2\mu/4} \quad (0 < \tau < 1).$$

Here the upper-tail exponent is $\mu((1 + \tau) \log(1 + \tau) - \tau)$; its second derivative in τ is $\mu/(1 + \tau) \geq \mu/2$, so it is at least $\mu\tau^2/4$. The lower-tail exponent is at least $\mu\tau^2/2$ by the same calculation. The union bound over at most Rp^3 symbol-progression pairs is at most $1/4$ when $n \geq L_{\text{bal}}$.

Now fix a nonconstant polynomial P of degree at most d on the whole field. It can take a value in S on at most dR points. At each such point the random word matches it with probability $1/R$,

independently. The number X of matches therefore has mean at most d . Its exponential moment is bounded by $\exp(d(e^u - 1))$. For any integer $t \geq d$, Markov's inequality, with $u = \log(t/d)$, yields

$$\mathbb{P}(X \geq t) \leq (ed/t)^t.$$

With $t = T \geq e^2 d$ the right-hand side is at most $e^{-T} \leq e^{-H}$. There are at most p^{d+1} coefficient lists, so the probability of any such global failure is at most $p^{d+1}e^{-H} = 1/4$. A bound for total agreement on the field also bounds agreement on every progression. With positive probability both (i) and (ii) hold.

On a progression satisfying (II.231.4), every symbol has at least $T + 1$ occurrences, whereas each nonconstant polynomial has fewer than T agreements. Repeat the proof of Theorem II.228.2, now charging a nonconstant polynomial T rather than dR . The optimum is the sum of the q largest symbol counts and strict rigidity follows. Finally, $R \leq p$ gives $\log(8Rp^3) = O(\log p)$, and $T = O((d + 1) \log p)$, proving the scale estimate. $\square$

Remark II.231.5 (A sharper implicit avoidance threshold). One may replace the convenient T in (II.231.3) by any integer $T \geq d$ satisfying

$$T \log \frac{T}{ed} \geq (d + 1) \log p + \log 4.$$

For fixed d its least solution is of order $\log p / \log \log p$. No union bound over polynomial lists is required: a uniform bound for each nonconstant graph suffices for the list envelope. The balanced-symbol scale may still dominate the final answer.

Corollary II.231.6 (A growing-alphabet regime). *Suppose $R = R(p) \leq p$, $d = d(p) < p$, and*

$$\frac{R(d + 1) \log p}{L_p} \longrightarrow 0, \quad L_p \leq p.$$

Then words can be chosen so that, simultaneously for every $q \leq R$ and every proper progression of length at least L_p , the optimum is exactly the mass of the q heaviest symbols and

$$\left| \frac{\text{Cov}_{q,d}(a; I)}{|I|} - \frac{q}{R} \right| = o(1).$$

The same asymptotic cover bound holds on adaptive unions of product boxes with those final-axis lengths, even for arbitrary dependence of the polynomial coefficients on the base point.

Proof. Choose $\tau = \tau_p \rightarrow 0$ sufficiently slowly that $R\tau_p^{-2} \log p = o(L_p)$. For example, for large p one may take $\tau_p = (R \log p / L_p)^{1/4}$. The hypothesis also ensures $RT = o(L_p)$ in Theorem II.231.4. The scale condition eventually holds. Each selected symbol mass lies within τ_p/R of $1/R$, so summing $q \leq R$ such masses creates error at most τ_p . The fibrewise proof of Theorem II.229.1 and summation over disjoint cells complete the argument. $\square$

II.232 The pinned Section 16 interface

Source 50, delivered Section 7.

This section translates only the definitions needed for the application. The source for these formulas is the pinned repository catalogue, including its corrections and proper-box requirements [218]. Bare references to “Lemma 16.10” in this section are not references to a theorem proved in this manuscript.

II.232.1 The multiple-multilinearity controls

For a dimension D , define

$$E_D = 2^{2^{D+8}}, \quad c(u, \gamma, D) = (\gamma u)^{E_D}, \quad q(u, \gamma, D) = (\gamma u)^{-E_D}, \quad (\text{II.232.1})$$

and

$$s(u, \gamma, D) = \left(\frac{2}{u\gamma} \right)^{2^{2^{D+6}}}. \quad (\text{II.232.2})$$

We use a lower-case $q(\cdot)$ only for this real-valued control function; an integer list length is written q_0 or Q in the application.

For a full-domain function $F : \mathbb{F}_p^D \rightarrow \mathbb{F}_p$, the property $\text{ML}_{\gamma,r}(F)$ in the catalogue says the following. For every $0 < \rho \leq 1$ and every proper common-step box P , there are a retained set $H \subseteq P$, a proper box partition (P_j) of P , and an integer q_0 such that

$$|H| \geq (1 - \rho)|P|, \quad (\text{II.232.3})$$

$$q_0 \leq q(\rho/r, \gamma, D)^r, \quad (\text{II.232.4})$$

$$\text{wid}(P_j) \geq \text{wid}(P)^{c(\rho/r, \gamma, D)^r}, \quad (\text{II.232.5})$$

and, on $H \cap P_j$, the function agrees pointwise with one of q_0 globally multiaffine polynomials. The choice of the matching polynomial can depend on the point. This is a cover by a union of graphs, not agreement with a single polynomial. The repository's relation-valued definition specializes to this form on a full graph.

Lemma II.232.1 (An identity-partition certificate). *Suppose $0 < \gamma \leq 1$, $r \geq 1$, and the image of F has size at most R . If*

$$R \leq (r/\gamma)^{E_{Dr}}, \quad (\text{II.232.6})$$

then $\text{ML}_{\gamma,r}(F)$ holds. In particular, every constant function has this property.

Proof. For every ρ , use the identity partition $P_1 = P$, retain all of P , and use the constant polynomials corresponding to the image of F . The graph budget is at least its value at $\rho = 1$, namely the right-hand side of (II.232.6). The required width exponent is $(\gamma\rho/r)^{E_{Dr}} \in (0, 1]$. For an integer width $w \geq 1$, $w^a \leq w$ when $0 < a \leq 1$; for $w = 0$, $0^a = 0$. Thus the width condition also holds. $\square$

II.232.2 The final-coordinate line-cover input

Let $k \geq 1$, so the ambient dimension is $k + 1$. For outer parameters $0 < \theta, \gamma \leq 1$, put

$$\theta_1 = (\theta\gamma/4)^{2^{k+5}}, \quad \delta = 2^{-37}(\theta_1/4)^{11/2}, \quad (\text{II.232.7})$$

$$\zeta = 2^{-s(\theta, \gamma, k)}, \quad t = \delta^{-2}s(\theta_1/8, \delta, k), \quad (\text{II.232.8})$$

$$r = (2^k - 1)\gamma^{-2}s(2^{-(k+2)}\theta, \gamma, k), \quad K = (k + 1)^2 2^{k+4}. \quad (\text{II.232.9})$$

For a loss $0 < \sigma \leq 1$, the two graph-count bounds are

$$Q_\gamma(\sigma) = q(\sigma/(2r), \gamma, k + 1)^r, \quad Q_\delta(\sigma) = q(\sigma/(2t), \delta, k)^t. \quad (\text{II.232.10})$$

Given an integer q_δ , the required line-cover width is

$$\ell(m, q_\delta) = \frac{\zeta}{2} m^{e(\sigma, q_\delta)}, \quad (\text{II.232.11})$$

$$e(\sigma, q_\delta) = \frac{c(\sigma/(2r), \gamma, k+1)^r c(\sigma/(2t), \delta, k)^t}{2K^{2^{k+1}q_\delta}}. \quad (\text{II.232.12})$$

All auxiliary parameters are positive; $\delta, \zeta \leq 1$ and $r, t \geq 1$ in the cases used below.

The predicate **Section16AllBoxLineCovers** requires, for every σ , every integer m , and every proper $(k+1)$ -box P with $m \leq \text{wid}(P)$, integers $q_\gamma \leq Q_\gamma(\sigma)$ and $q_\delta \leq Q_\delta(\sigma)$, together with a proper partition $S_u = T_u \times I_u$ of P and a set $E \subseteq P$ of mass at least $(1-\sigma)|P|$. The cells have width at least (II.232.11); on each final-coordinate fibre in E , the function is covered by q_γ globally affine functions of the final coordinate. Their coefficients may depend on the cell and on the base point.

Lemma II.232.2 (Identity line covers). *If F takes at most R values and*

$$R \leq Q_\gamma(1), \quad Q_\delta(1) \geq 1,$$

*with the parameters satisfying $r, t \geq 1$, $0 < \gamma, \delta, \zeta \leq 1$, then F satisfies **Section16AllBoxLineCovers**.*

Proof. Use $q_\gamma = R$, $q_\delta = 1$, the identity partition, $E = P$, and the R constant final-coordinate affine functions. The bounds in (II.232.10) increase when σ decreases, so the graph budgets hold for every σ . Both factors in the numerator of (II.232.12) lie in $(0, 1]$, and its denominator is at least 2. Thus $0 < e \leq 1$. For $m \geq 1$, $\ell(m, 1) \leq m \leq \text{wid}(P)$; for $m = 0$ the width requirement is zero. The identity box has the required base/final-axis product structure and properness. All conditions are therefore witnessed explicitly. $\square$

II.232.3 The precise two-input implication

The catalogue's **lemma_16_10** asserts, for a certain named full or partial map ϕ_1 on a domain B_1 , that

$$\begin{aligned} &[\text{every fixed-final-coordinate section is } \text{ML}_{\gamma, r_0}] \quad \text{and} \quad [\text{all the line covers above exist}] \\ &\implies \text{ML}_{\gamma, 1}(\phi_1), \quad (\text{II.232.13}) \\ &r_0 = \gamma^{-2} s(2^{-(k+2)} \theta, \gamma, k). \end{aligned}$$

The auxiliary names B_1 and ϕ_1 are not unrestricted variables in the source, so a counterexample must also realize them using the source's cube selection construction. That realization is included in the next section.

II.233 A counterexample to the two-input form of Lemma 16.10

Source 50, delivered Section 8.

Theorem II.233.1 (Fixed-parameter, arbitrarily large prime counterexamples). *Let*

$$k = 1, \quad \theta = \gamma = 1, \quad E = E_2 = 2^{1024}, \quad Q = 2^E, \quad R = 4Q, \quad c = Q^{-1}. \quad (\text{II.233.1})$$

For every prime

$$p > 2^{4096Q^4}, \quad (\text{II.233.2})$$

there are witnesses to both premises of the pinned `lemma_16_10`, with $B_1 = \mathbb{F}_p^2$, for which its conclusion is false. More specifically, they can be chosen so that $\phi_1(h, y) = a(y)$ for an R -valued word and every admissible target cover at loss $\rho = 1/2$ covers at most $5p^2/16$ points of the full box.

Proof. We verify the assertion in five steps.

1. Construct the word. Choose $S = \{0, 1, \dots, R-1\} \subset \mathbb{F}_p$ and put $\varepsilon = 1/(4R)$. Since $\log 2 \geq 1/2$, condition (II.233.2) gives

$$\log p > 2048Q^4 = \frac{128R^2}{c^2}.$$

Lemma II.230.4 implies $p \geq R$ and

$$p^c \geq L(p, R, \varepsilon), \quad p^c \geq R^2. \quad (\text{II.233.3})$$

Use Theorem II.230.2 with the uniform distribution on S to choose $a : \mathbb{F}_p \rightarrow S$ balanced within ε on every proper progression of length at least L . For a fully specified deterministic choice, order $\mathbb{F}_p$ and S by their integer representatives and use the least-score, least-symbol tie-breaking rule of Theorem II.230.3. No value of this enormous word needs to be computed to establish its existence or the properties proved here.

2. Realize the source's named map and domain. Take $B = \mathbb{F}_p^2$, $H_1 = \mathbb{F}_p$, the anchor $x_0 = 0$, and

$$\phi(h, y) = a(y).$$

For every sidelength h , let Y_h consist of *all* cube elements in `Section16CubeElement B h`. In dimension one such an element is a horizontal pair of vertices with sidelength h , together with its final coordinate. Because B is full, every such pair is admitted, including the pair based at 0 with any final coordinate y .

The source predicate `Section16GoodInducedPair` requires membership of h in H_1 and an element of Y_h whose base is x_0 and whose final coordinate is y . These requirements now hold for every (h, y) . Consequently

$$B_1 = \text{section16GoodDomain}(B, H_1, Y, x_0) = \mathbb{F}_p^2.$$

The map `section16PhiOne` is $\phi(x_0 + h, y)$, which here equals $a(y)$. Thus the named domain and map are genuinely realized, not substituted by an informal surrogate.

3. Verify the cross-section premise. At the chosen parameters,

$$r_0 = s(1/8, 1, 1) = 16^{2^{128}} \geq 2. \quad (\text{II.233.4})$$

For each fixed y , the map $h \mapsto a(y)$ is constant on all of $\mathbb{F}_p$. Lemma II.232.1 gives ML_{1, r_0} for that section, using one graph and retaining every point on every proper box. This proves exactly `FinalCoordinateSectionsMultiplyLinear` in the antecedent.

4. Verify every line-cover requirement. Since $k = 1$, the r in (II.232.9) equals r_0 . The auxiliary values in (II.232.7)–(II.232.9) satisfy $0 < \delta, \zeta \leq 1$ and $t \geq 1$. For every $0 < \sigma \leq 1$,

$$Q_\gamma(\sigma) = \left(\frac{2r_0}{\sigma}\right)^{Er_0} \geq 2^{2E} = Q^2 \geq 4Q = R,$$

$$Q_\delta(\sigma) = \left(\frac{2t}{\delta\sigma}\right)^{E_1t} \geq 1.$$

Here $Q \geq 4$ and $r_0 \geq 2$ justify the comparisons. Set $q_\gamma = R$ and $q_\delta = 1$. Lemma II.232.2 supplies all boxes, all losses, and all integer scales m , with the identity partition and no exceptional points. Therefore `Section16AllBoxLineCovers` holds.

5. Refute the conclusion. Apply the asserted conclusion $\text{ML}_{1,1}(\phi_1)$ to the full proper box $P = \mathbb{F}_p^2$ represented by two length- p , step-one axes, at loss $\rho = 1/2$. Its graph and width requirements become

$$q_0 \leq q(1/2, 1, 2) = 2^E = Q, \quad \text{wid}(P_j) \geq p^{c(1/2, 1, 2)} = p^c. \quad (\text{II.233.5})$$

In each nonempty cell $P_j = T_j \times I_j$, the final axis is a proper modular progression of length at least p^c . By (II.233.3) it is long enough for the balance property and has length at least R^2 . Multiaffine graphs restrict to degree-at-most-one polynomials in the final coordinate. Corollary II.229.2 therefore bounds the covered mass, after *any* such partition, by

$$q_0 \left(\frac{1}{R} + \frac{1}{4R}\right) |P| \leq \frac{5Q}{4R} p^2 = \frac{5}{16} p^2.$$

Every retained point has to be covered because B_1 is the entire domain. But (II.232.3) demands $|H| \geq p^2/2$. Since $5/16 < 1/2$, this is impossible. Both premises hold and the conclusion fails. $\square$

[write] **7 October 2026; sources 47–53.** Checked at the write: $E_2 = 2^{2^{10}} = 2^{1024}$, $r_0 = s(1/8, 1, 1) = 16^{2^{128}}$, $5Q/(4R) = 5/16 < 1/2$ for $R = 4Q$, and $\log p > 2048Q^4 = 128R^2/c^2$ when $p > 2^{4096Q^4}$ (exact exponent arithmetic). The negation of `lemma_16_10` is now a theorem of the repository (note in Section F.55); this is a counterexample to the packaged two-premise abstraction, not to Gowers’s Lemma 16.10 with its upstream hypotheses, as source 50 says.

[write] **7 October 2026; sources 60–69.** *The catalogue entry has since been repaired* (456b6e3ad, 7 October 2026, 12:20; checked at HEAD fb9602e55). The predicate this theorem refutes is kept, unchanged in content, as `lemma_16_10_printed_unit_encoding` (`Section16.lean`, line 852), and `lemma_16_10_packaged_premises_counterexample` (`Proofs16PackagedPremisesCounterexample.lean`, line 14) now negates it under that name. The name `lemma_16_10` (line 912) now denotes a repaired statement, the cover the printed proof establishes from the same two premises (graph count $\max(b, \binom{r}{2}b^2)$, exceptional mass $\rho|P|$), proved by `lemma_16_10_holds` (`Proofs16Lemma10.lean`, line 12) through the positive interface named in this source’s front matter; `FORMALIZATION_STATUS.txt` counts it (114 exact companions, 6 open statements) and states that it is not yet kernel-checked. This theorem is unaffected: it refutes the printed unit-parameter form, which the repair does not claim. Earlier mentions of `lemma_16_10` as open or false in this report mean the renamed predicate (Section V.1).

II.233.1 Why the usual possible repairs do not address this example

The example uses prime moduli, proper boxes, distinct residues, a full domain, positive width exponents, and the allowed parameter values $\theta = \gamma = 1$. It therefore does not rely on composite-modulus torsion, zero iteration parameters, inflated formal widths, sparse domains, or the limit

$\theta \rightarrow 0$. There are counterexamples for every prime beyond the explicit bound, so introducing an arbitrary threshold $p \geq p_0(\theta, \gamma, k)$ does not repair this two-input implication. The fibre cover is exact, so improving its exceptional-set estimate cannot repair it either.

The obstruction is to a particular graph budget. It does not say that the function lacks any multiple-multilinear description: the R constant graphs already provide one. What fails is replacing the permitted input branch count by the smaller count in (II.233.5) while retaining half the mass on long cells.

Corollary II.233.2 (Hamming robustness). *Let a be a word from Theorem II.233.1. If $a' : \mathbb{F}_p \rightarrow S$ differs from a on fewer than $3p/16$ coordinates, then the same two-premise counterexample is obtained from $F'(h, y) = a'(y)$.*

Proof. The premises are unchanged: the slices remain constant and the same R constant fibre graphs cover everything. For any proposed target partition and lists, replacing a by a' changes agreement at at most $p|\{y : a(y) \neq a'(y)\}| < 3p^2/16$ points. The universal $5p^2/16$ upper bound for a therefore remains strictly below $p^2/2$ for a' . $\square$

Proposition II.233.3 (Both orientations of coordinate sections). *In Theorem II.233.1, every one-dimensional coordinate section, in either orientation, is ML_{1, r_0} with r_0 as in (II.233.4). Hence adding that extra coordinate-section hypothesis alone does not repair the assertion.*

Proof. Sections with the final coordinate fixed are constant. Sections with the base coordinate fixed equal the R -valued word a . For those, apply Lemma II.232.1 in dimension one. Indeed $r_0 = 2^{4 \cdot 2^{128}} \geq 2E$, so

$$r_0^{E_1 r_0} \geq 2^{r_0} \geq 2^{2E} = Q^2 \geq R.$$

This verifies the necessary identity-partition graph budget. $\square$

Proposition II.233.4 (Invisibility to horizontal differences). *For every word a and $F(h, y) = a(y)$, every positive-order alternating cube difference in the base variables vanishes identically. In the one-dimensional base used above, the source's induced cube map is identically zero.*

Proof. All vertices of a base-variable cube have the same final coordinate, so all values of F are equal. The sum of the alternating signs is $(1 - 1)^k = 0$ for a positive-dimensional cube. In dimension one the value is simply $a(y) - a(y) = 0$. $\square$

This last observation explains why some additional compatibility tests may still miss the problem. Any proposed repair must account for one-coordinate summands that disappear under the differences it tests. This does not imply the full upstream product property. In fact, Proposition II.234.1 proves that the witnesses fail that property at $\gamma = 1$.

II.233.2 A parameter-general version

Proposition II.233.5 (The obstruction is not restricted to endpoint parameters). *Fix $k \geq 1$ and $0 < \theta, \gamma \leq 1$. Put*

$$E = E_{k+1}, \quad Q = (2/\gamma)^E, \quad R = \lceil 4Q \rceil, \quad c = Q^{-1}.$$

For all sufficiently large primes the same full-domain construction satisfies the two premises at these parameters and violates the conclusion at loss $1/2$.

Proof. The cross-sections are constant, with the prescribed $r_0 \geq 1$. The line-cover iteration $r = (2^k - 1)r_0$ is at least two. Therefore

$$Q_\gamma(1) = (2r/\gamma)^{Er} \geq (2/\gamma)^{2E} = Q^2.$$

Since $Q \geq 5$ and $R \leq 4Q + 1 \leq 5Q$, we have $Q_\gamma(1) \geq R$. The delta-side budget is at least one, so the identity line cover works. Choose p large enough for Lemma II.230.4, use the uniform word on R symbols, and realize the full good domain by taking all cubes and anchor zero. The target graph count is at most Q and the target width exponent is c . Hence the covered fraction is at most $5Q/(4R) \leq 5/16$, again less than $1/2$. All parameters other than the prime and the word are fixed in this argument. $\square$

II.234 A repair interface that retains complexity

Source 50, delivered Section 9.

II.234.1 What a valid replacement must preserve

There are two separate tasks in the Section 16 lifting argument. One is to turn fibrewise affine data and slice information into multiaffine graph covers. The other is to fit the resulting count and width into prescribed control functions. Our result is a barrier to the second task when only the bundled inputs are retained; it is not a barrier to interpolation itself.

A replacement theorem can keep its output parameters explicitly dependent on the input branch count, use a suitably larger iteration parameter, require extra structural hypotheses, or weaken the retained mass. Merely retaining a formal assertion that the parameters can later be absorbed is not a repair. Corollary II.231.2 gives a necessary check: on uniform R -branch words, an output list retaining $1 - \eta$ of long boxes must have at least $(1 - \eta - o(1))R$ branches.

II.234.2 The full-domain endpoint product property

The upstream product property is not merely an unspecified possible source of additional information. At its endpoint, it can be characterized exactly. This also verifies a concrete limit on what the counterexample disproves.

For a full-domain $G : \mathbb{F}_p^D \rightarrow \mathbb{F}_p$, the repository's product property at parameter γ requires the following for every finite family of parallel coordinate restrictions ψ_i , every set $A \subseteq \mathbb{F}_p$, and nonnegative weights w on A :

$$\sum_{\substack{x_1+x_2=x_3+x_4 \\ x_t \in A \\ \psi_i(x_1)+\psi_i(x_2)=\psi_i(x_3)+\psi_i(x_4) \ \forall i}} \prod_{t=1}^4 w(x_t) \geq \gamma^{8m} p^{-1} \left(\sum_{x \in A} w(x) \right)^4, \quad (\text{II.234.1})$$

where m is the number of restrictions [221].

Proposition II.234.1 (Endpoint characterization). *For a full-domain map $G : \mathbb{F}_p^D \rightarrow \mathbb{F}_p$, $D \geq 1$, the product property with $\gamma = 1$ is equivalent to G being globally multiaffine. Consequently it implies $\text{ML}_{1,1}(G)$ with one graph. The witnesses in Theorem II.233.1 do not have this upstream endpoint product property.*

Proof. Assume the product property at $\gamma = 1$. Take one coordinate restriction ψ , $A = \mathbb{F}_p$, and $w = 1$. There are exactly p^3 ordered additive quadruples, and (II.234.1) requires all p^3 of them to be respected. In particular,

$$\psi(x) + \psi(y) = \psi(x + y) + \psi(0) \quad (x, y \in \mathbb{F}_p).$$

Thus $\psi - \psi(0)$ is additive. Since the additive group of $\mathbb{F}_p$ is generated by 1, it equals $x \mapsto sx$ for some $s \in \mathbb{F}_p$. Every coordinate restriction is therefore affine.

A separately affine function on a full product of fields is multiaffine. To verify this without a separate theorem, induct on D and write

$$G(h, y) = G(h, 0) + y(G(h, 1) - G(h, 0)).$$

The two base functions are separately affine, so the induction hypothesis expresses them as multiaffine polynomials. The displayed formula gives the joint multiaffine representation.

Conversely, for a multiaffine G , every coordinate restriction is affine and therefore respects every additive quadruple. For any finite family of them, the sum on the left of (II.234.1) is the ordinary weighted additive energy. Extend w by zero outside A and let $r(s) = \sum_x w(x)w(s - x)$. Cauchy–Schwarz gives

$$\sum_s r(s)^2 \geq p^{-1} \left(\sum_s r(s) \right)^2 = p^{-1} \left(\sum_x w(x) \right)^4.$$

This is exactly the product-property bound at $\gamma = 1$. The identity partition with the single graph G proves $\text{ML}_{1,1}(G)$.

Finally, our balanced word has every one of its R symbols present, since its full-field symbol masses are at least $3/(4R)$. Here $2 \leq R < p$. A nonconstant affine map on $\mathbb{F}_p$ has image size p , and a constant one has image size one. Thus a is not affine. The coordinate restriction $y \mapsto a(y)$ rules out the endpoint product property. $\square$

[write] **7 October 2026; sources 47–53.** *Formalized* (checked at HEAD 06b30ca52): `section16_unit_product_iff_multilinear` (`Proofs16MultilinearProduct.lean`, line 89) proves that over a prime field the full-domain product property at $\gamma = 1$ is equivalent to multilinearity, with `HasProductProperty.unit_multilinear` (`Proofs16ProductEndpoint.lean`, line 146) for the forward direction; `FORMALIZATION_STATUS.txt`: it “supplies the positive endpoint requested by source 50’s plan”.

The equivalence is an elementary endpoint fact, not a claim to a new inverse theorem. It pinpoints a genuine hypothesis lost by the two-input abstraction. It does not establish a corrected lifting statement for $0 < \gamma < 1$, nor does it show that the full product property is inherited by every intermediate object at the same parameter. Those are different quantitative questions.

II.234.3 The existing positive local lift

The pinned repository already contains a theorem `section16_local_affine_lift` with explicit output data [222]. Its role is important here because it avoids the false parameter reset. In a simplified mathematical reading, start with an affine fibre cover of size $q_{\text{in}} > 0$, a slice parameter $s \geq 1$, losses $\tau, \epsilon \in (0, 1]$, and an integer sampling parameter r_{sam} satisfying

$$6q_{\text{in}} \leq r_{\text{sam}}\tau^2.$$

For a base scale $m \geq 4$ and an integer $v \geq 1$, the theorem assumes

$$v^2 + 1 \leq (m/8)^{c(\epsilon/(r_{\text{sam}}s), \gamma, k)^{r_{\text{sam}}s}}.$$

It produces an integer p_0 satisfying

$$p_0 \leq [q(\epsilon/(r_{\text{sam}}s), \gamma, k)]^{r_{\text{sam}}s}. \quad (\text{II.234.2})$$

The graph-index type in that theorem has cardinality

$$r_{\text{sam}}p_0 + r_{\text{sam}}^2p_0^2. \quad (\text{II.234.3})$$

The output box width is at least $v - 1$, and the retained fraction is at least $1 - 2\tau - \epsilon$. For a partial domain, the coverage conclusion applies on that domain intersected with the retained set. In a full-domain application it is a cover of the entire retained set.

We cite this as a source theorem, not as a newly verified Lean result: no kernel rebuild of the repository was performed for this article. The important point is that (II.234.3) continues to depend on the sampling budget and therefore on q_{in} . Its statement does not silently identify that count with $q(\rho, \gamma, k + 1)$.

[write] 7 October 2026; sources 47–53. *Changed since the pin* (checked at HEAD 06b30ca52). `section16_local_affine_lift` (`Proofs16LocalAffineLift.lean`, line 13) now assumes the linear sampling budget $6q \leq r\tau$ in place of $6q \leq r\tau^2$ (93e36c6df) and indexes its graphs by `section16CompressedCandidateCount` $rp = \max\{p, \binom{r}{2}p^2\}$ in place of $rp + r^2p^2$ (081abc4d3); the description above is that of the pin. The count still depends on the sampling budget, so the point made above stands.

II.234.4 An elementary label-preserving interpolation lemma

The next proposition records a transparent sufficient condition for gluing. It is standard interpolation, not a new substitute for the source theorem. Its purpose is to exhibit an interface in which the branch budget is visible.

Proposition II.234.2 (Label-preserving two-anchor gluing). *Let $u \neq v$ in $\mathbb{F}_p$. Suppose F is defined on a subset D of a product $T \times I$. For each branch $j \in \{1, \dots, R\}$, assume that its final-variable map $\ell_{j,h}(y)$ is affine in y . On a common base-box partition, suppose its anchor values $\ell_{j,h}(u)$ and $\ell_{j,h}(v)$ are covered by respective lists of A_j and B_j multiaffine functions of h . Outside a base exceptional set of relative mass β , assume these covers and the branch representation of F hold. Then a multiaffine cover on the corresponding product cells exists with at most*

$$\sum_{j=1}^R A_j B_j \quad (\text{II.234.4})$$

graphs, losing at most the base exceptional fraction β there. A separate pointwise representation failure of mass α increases the total loss by at most α .

Proof. For every candidate pair (U, V) of anchor functions for branch j , form

$$L_{j,U,V}(h, y) = \frac{v - y}{v - u} U(h) + \frac{y - u}{v - u} V(h).$$

This is multiaffine in h and affine in y , hence multiaffine jointly. At any good h , a correct pair of candidates agrees with the two anchor values, so its interpolant is exactly $\ell_{j,h}$ for all y . There are $A_j B_j$ pairs for branch j . Take their union over j . The exceptional-set assertion follows by multiplying the base exceptional set by the final axis and adding any separate pointwise failures. $\square$

For the model words every branch has $A_j = B_j = 1$, so this construction uses R constant graphs. Theorem II.231.1 shows why a general cover cannot replace that number by an arbitrary smaller budget while keeping almost all the mass. A degree- d version uses $d + 1$ distinct anchors and $\sum_j \prod_{t=0}^d A_{j,t}$ interpolants; it follows from the same Lagrange formula and preserves degree at most d in the final variable.

II.234.5 Formal interface recommendations

Keep a local lifting theorem parameterized by the actual input count, sampling budget, slice control, retained mass, and output width. Make any subsequent comparison with a standard control function a separate theorem, with all outer parameters universally quantified in its statement. The counterexample in Section II.233 should be a negative test for proposed comparisons. No theorem label should receive a proved status merely because a nearby interpolation theorem has been checked.

A more ambitious repair may strengthen the upstream structural assumptions. That requires determining which of those assumptions really constrain the one-coordinate summands in Proposition II.233.4; it cannot be inferred from the existence of the two bundled covers alone.

II.235 An acyclic replacement for the Section 16 cover theorem

Source 49, delivered Section 8.

This section proves a quantitative implication from seven finite interfaces appearing before the closing numerical comparison in Section 16 of Gowers's paper [1, 210]. The earlier project audit did not establish that comparison for the displayed controls. Here the term *acyclic* refers to explicit replacement recurrences in which every recursive call strictly decreases the active dimension. The retained global set is chosen before the local error tolerance. Graph counts, width exponents, minimum widths, and prime thresholds are kept separate. The result supplies a replacement for the logical conclusion of Corollary 16.11. It does not establish the paper's printed control functions or a new global bound for arithmetic-progression-free sets.

The sampling and shared-interpolation ingredients have already been developed in the earlier interpolation report [215]. They are restated here to make the induction complete. The additional work is the full recurrence, its quantifier order, and its treatment of unused coordinates. The latter point is useful: a partition in fewer coordinates cannot simply be assumed to extend to a partition whose common difference is synchronized in all coordinates.

II.235.1 The numerical comparison being replaced

For clarity, we reproduce the earlier project's algebraic audit [215]. This calculation concerns one displayed numerical comparison, not the truth of the source's structural assertion. Put

$$\mathcal{A}_j = 2^{2^{j+8}}, \quad \mathcal{Q}(t, \gamma, j) = (\gamma t)^{-\mathcal{A}_j}, \quad \mathcal{B}_k = 2^{2^{k+6}}.$$

With $0 < \theta, \gamma, \rho \leq 1$ and a sample size $r \geq 1$, the auxiliary quantities in the published closing comparison have the form

$$\mathcal{R} = \gamma^{-2} \left(\frac{2^{k+3}}{\theta \gamma} \right)^{\mathcal{B}_k}, \quad \Omega = 4r^2 \mathcal{R}.$$

The comparison on printed page 575 is

$$\mathcal{Q}(\rho/(4\Omega), \gamma, k)^\Omega \leq \mathcal{Q}(\rho, \gamma, k+1). \quad (\text{II.235.1})$$

Writing $b = \log(1/(\gamma\rho)) \geq 0$ and using $\mathcal{A}_{k+1} = \mathcal{A}_k^2$, it is equivalent to

$$\Omega\{b + \log(4\Omega)\} \leq \mathcal{A}_k b.$$

But $\mathcal{R} \geq 2^{(k+3)\mathcal{B}_k} > \mathcal{A}_k$ for every $k \geq 1$, and hence $\Omega > \mathcal{A}_k$. The inequality fails. The edited transcription has a different, same-dimension version on the left, which fails still more directly. The two versions should not be identified with one another.

[write] **7 October 2026; sources 47–53.** This is source 19’s audit of the closing comparison (Section V.80.2; formally `section16_same_dimension_closing_comparison_fails`, `Proofs16ClosingComparison.lean`, line 25, at HEAD 06b30ca52), as source 49 says. Source 50, of the same day, shows that the two-input catalogue statement `lemma_16_10` is itself false, not only the comparison (Section II.226); this section’s theorem is a conditional replacement with new controls, from the stronger hypotheses II–I7.

The smaller anchor budget from [215] is useful, but by itself does not justify (II.235.1). The remedy below defines new controls from the operations actually performed. It permits a weaker width exponent and carries a separate minimum width, instead of trying to absorb the entire refinement into these particular printed powers.

II.235.2 Active coordinates and the product property

Throughout this section, p is prime, Fourier transforms on $\mathbb{F}_p$ are unnormalized, and $|t|_p$ denotes the least absolute residue of $t \in \mathbb{F}_p$. Thus

$$\widehat{g}(r) = \sum_{x \in \mathbb{F}_p} g(x) \exp(-2\pi i r x / p), \quad \text{Bohr}(K, \rho) = \{t : |rt|_p \leq \rho p \text{ for all } r \in K\}.$$

A box is a product of nonempty injective arithmetic progressions having one common nonzero difference. Its width is the smallest side length. A multilinear polynomial in this section means a polynomial affine in each variable separately, so it includes all lower-degree terms.

Definition II.235.1 (Active-coordinate cover). Let $\Lambda \subseteq \mathbb{F}_p^j \times \mathbb{F}_p$, let $D \geq j$, and choose j distinct ambient coordinates, in a specified order, with independent fixed translations. Write $\pi : \mathbb{F}_p^D \rightarrow \mathbb{F}_p^j$ for the resulting projection and translation map. On a box $P \subseteq \mathbb{F}_p^D$, a cover with parameters $(Q, c, T; \varepsilon)$ consists of a set $E \subseteq P$, $|E| \geq (1 - \varepsilon)|P|$, and a partition of P into boxes of width at least m^c , where m is the width of P and $m \geq T$. On each cell there are at most Q polynomials depending only on the active coordinates, multilinear in those coordinates, whose values include every y with $(\pi(z), y) \in \Lambda$ whenever $z \in E$ belongs to that cell.

The exceptional set need not depend only on the active coordinates. The models themselves must do so. Covers are preserved by restricting Λ , by translating active coordinates, and by permuting them.

Definition II.235.2 (Product property). A partial map $\phi : A \subseteq \mathbb{F}_p^j \rightarrow \mathbb{F}_p$ has the γ -product property if the following holds. Choose any coordinate direction, any $s \geq 1$ restrictions $\psi_1, \dots, \psi_s$ to parallel coordinate lines, any common domain $U \subseteq \mathbb{F}_p$, and any weight $w : U \rightarrow [0, \infty)$. The sum of $w(a)w(b)w(c)w(d)$ over quadruples with

$$a + b = c + d, \quad \psi_i(a) + \psi_i(b) = \psi_i(c) + \psi_i(d) \quad (1 \leq i \leq s)$$

is at least

$$\gamma^{8s} p^{-1} \left(\sum_{x \in U} w(x) \right)^4.$$

A relation $\Gamma \subseteq \mathbb{F}_p^j \times \mathbb{F}_p$ has that property if every partial map with graph contained in Γ does.

This is the paper's product property, with the ambient prime denoted by p . It is inherited by subrelations and proper coordinate sections.

II.235.3 The finite interfaces used as hypotheses

The implication below assumes the following corrected interfaces. Their proofs are earlier parts of the source development; reproducing all of them would amount to reproducing much of the paper. Stating the interfaces explicitly identifies the dependency boundary of the new induction. None invokes the conclusion being proved here.

11. Order-eight extraction. If a map on a set of density b has at least $\rho(bp)^3$ additive quadruples, then it has a restriction to an order-eight Freiman map on at least

$$2^{-1882} \rho^{1164} bp \quad (\text{II.235.2})$$

points.

12. Bohr linearization of an exact Freiman map. An order-eight Freiman map on a set of density b admits a set K of at most $16b^{-2}$ frequencies such that its restriction to any progression of length at most L is affine whenever its difference d satisfies

$$d \in \text{Bohr}(K, b/(32\pi L)). \quad (\text{II.235.3})$$

The statement concerns the original domain's global density. No local density assumption in the progression is required.

13. Many respected arrangements. A γ -product map on a set of density b in $\mathbb{F}_p^{k+1}$, $k \geq 1$, respects at least

$$b^{7 \cdot 4^{k+1}} \gamma^{21k \cdot 4^{k+1}} p^{17k+15} \quad (\text{II.235.4})$$

eight-arrangements. Here an eight-arrangement is an ordered family of sixteen coordinate cubes with common side vector $h \in \mathbb{F}_p^k$, the i th cube lying in last-coordinate section x_i , where $x_1 + \cdots + x_8 = x_9 + \cdots + x_{16}$. The arrangement is respected if the corresponding signed cube sums of the map satisfy the same equality. Repetitions are counted in this convention.

14. Riesz selection. Set

$$H_k = 2^{2^{k+4} + k + 3}. \quad (\text{II.235.5})$$

If the respected count is at least $\alpha b^{15} p^{17k+15}$ and $0 < \eta_0 \leq 1$, there is a subset of the original map domain with at least

$$(\alpha \eta_0 / 4)^{H_k} b^{15} p^{17k+15} \quad (\text{II.235.6})$$

arrangements, at least a proportion $1 - \eta_0$ respected, provided

$$(\alpha \eta_0 / 4)^{H_k} b^{15} p \geq k 3^{2^{k+4}}. \quad (\text{II.235.7})$$

I5. A Bohr model for an almost homomorphism on a domain. Let a finite domain X have a projection $r : X \rightarrow \mathbb{F}_p$, with fibres of size at most M and $|X| = bMp$. Suppose a map on X respects at least a proportion $1 - 2^{-43}$ of its order-eight additive relations. Put $g(x) = |r^{-1}(x)|$ and

$$\begin{aligned} K_b &= \{t : |\widehat{g}(t)| \geq 2^{-37} b^{11/2} Mp\}, \\ R_b &= 2^{74} b^{-10}, \quad z_b = 2^{-155 R_b} b^{18 R_b} / R_b. \end{aligned} \tag{II.235.8}$$

Then $|K_b| \leq R_b$, and there is $Y \subseteq X$ with $|Y| \geq b^6 |X| / 20000$ such that the map restricted to Y is a $\text{Bohr}(K_b, z_b)$ -homomorphism. In particular its value is constant on each fibre of Y , and the resulting partial function of the projected coordinate is affine on any progression of length at most L with difference in $\text{Bohr}(K_b, z_b/L)$.

I6. Product properties of large spectra. For $f : \mathbb{F}_p^{k+1} \rightarrow \mathbb{C}$ with $|f| \leq 1$, define

$$f_h(x) = \sum_{z \in \mathbb{F}_p^k} \prod_{\omega \in \{0,1\}^k} \mathcal{C}^{|\omega|+k} f(z + \omega h, x),$$

where $\mathcal{C}$ is complex conjugation and multiplication ωh is coordinatewise. The relation

$$\{(h, t) : |\widehat{f_h}(t)| \geq \delta p^{k+1}\} \tag{II.235.9}$$

has the δ -product property. Also, for $f : \mathbb{F}_p \rightarrow \mathbb{C}$, $|f| \leq 1$, a partial graph selecting a Fourier coefficient of modulus at least γp from each k -fold multiplicative derivative has the γ -product property. The elementary large-coefficient consequence of the paper's Lemma 3.1 is used only in Corollary II.235.11 below.

I7. Simultaneous multilinear phase partition. For $D \geq 2$ and an integer $q \geq 1$, define

$$\begin{aligned} K_D &= D^2 2^{D+3}, \quad a_D(q) = K_D^{-2^D q}, \\ L_D(q) &= 2^{\left(2^{40D^3+1}+1\right) K_D^{2^D q}}. \end{aligned} \tag{II.235.10}$$

Given q multilinear polynomials on a D -box of width $w \geq L_D(q)$, there is a partition into boxes of width at least $w^{a_D(q)}$ on each of which every polynomial has image diameter at most $2pw^{-a_D(q)}$ in the least-residue metric.

Interface	Corrected source declaration	Transcript line(s)
I1	Corollary 7.6	1336
I2	Lemma 7.8; Corollary 7.9	1350; 1389
I3	Lemma 14.8	3121
I4	Lemma 15.5	3223
I5	Theorem 10.13; Corollary 10.14	2219; 2253
I6	Lemmas 14.2 and 14.3	2969; 2980
I7	Corollary 5.11	978
Large-coefficient step	Lemma 3.1	372

Table II.235.1: Dependency crosswalk in the pinned `sz-thm-gowers-proof.tex` transcript [210]. These are source-declaration locations, not new Lean verification claims.

[write] 7 October 2026; sources 47–53. Checked at HEAD 06b30ca52. The line numbers are those of the transcription at the pin 8c12c906e; at HEAD three Section 13 commits have added 21 lines, so Lemmas 14.2, 14.3, 14.8 and 15.5 now stand at lines 2990, 3001, 3142 and 3244, and the other rows are unchanged. All eleven cited statements have exact companions in the ledger at HEAD (corollary_7_6, lemma_7_8, corollary_7_9, lemma_14_8, lemma_15_5, theorem_10_13, corollary_10_14, lemma_14_2, lemma_14_3, corollary_5_11, lemma_3_1). Whether each formal statement implies the interface I1–I7 exactly as formulated above was not checked at the write; source 49 does not certify them either, and its theorem stays conditional on I1–I7.

II.235.4 Two elementary cover operations

Lemma II.235.3 (Sequential refinement). *Suppose s relations or partial functions on one ambient box have cover triples (Q_i, c_i, T_i) at error ε/s . Put $C = \prod_{i=1}^s c_i$. A simultaneous refinement has width exponent C , total error at most ε , and a sufficient initial width threshold is $\max_i T_i^{1/C}$. A union needs at most $\sum_i Q_i$ graph models, and a sum of functions on a common domain needs at most $\prod_i Q_i$ models.*

Proof. Apply the cover for the first relation, then apply the second cover inside every resulting box, and continue. At stage i the width is at least $m^{c_1 \cdots c_i}$, and all previous covers remain valid on subboxes. Since $c_i \leq 1$, every input width is at least m^C , which exceeds each required T_i . The discarded masses add, because the boxes at any stage partition the original box. Lists concatenate for a union; all choices of one model from each list give the models for a sum. Models retain the union of their active coordinate sets. $\square$

Lemma II.235.4 (Anchors, including singleton classes). *Let $S = T \times I$ be a nonempty box with distinguished coordinate $x \in I$. Suppose that each row of a partial function on S is partitioned into at most q classes, with the function affine on each class. For every integer $r \geq 1$, one can choose at most r distinct anchors in I so that, except for at most $2q|S|/(r+1)$ points, each point is either an anchor point or belongs to a class containing two distinct anchors. If every chosen anchor slice has at most v multilinear models after a common refinement, then these retained points are covered by at most*

$$rv + \binom{r}{2} v^2 \tag{II.235.11}$$

multilinear models, including the distinguished coordinate.

Proof. If $|I| \leq r$, choose every point of I and use the constant extensions of the anchor slices; there is no loss. Otherwise sample an r -subset uniformly. Choose a test point uniformly outside this subset. For a fixed class C , let Z count its points in the resulting uniform $(r+1)$ -subset. Exchangeability gives

$$\Pr(\text{test point is in } C, \text{ at most one anchor is in } C) = \mathbb{E} \left[\frac{Z}{r+1} \mathbf{1}_{Z \leq 2} \right] \leq \frac{2}{r+1}.$$

Sum over at most q classes and over all rows. The nonanchor fraction of the box is at most one, and all anchor points are retained. Averaging therefore gives the required anchor set.

Write two distinct anchors as x_i, x_j . If $u_i(h), u_j(h)$ are models for their slices, their interpolant is

$$\frac{(x_j - x)u_i(h) + (x - x_i)u_j(h)}{x_j - x_i}. \tag{II.235.12}$$

It is multilinear in the active h coordinates and affine in x . The denominator is nonzero because the anchors are distinct in a prime field. Include all v^2 interpolants for each anchor pair and all v constant extensions for each anchor. On a class hit twice, the appropriate interpolant agrees with the class's affine map. An anchor point is covered by its slice's constant extension. This also explains why singleton classes cause no exception to the proof. $\square$

II.235.5 The replacement theorem and its base case

Theorem II.235.5 (Acyclic cover theorem from I1–I7). *Assume interfaces I1–I7. The explicit recurrences below define, for each $j \geq 1$ and $0 < a, \gamma \leq 1$, an integer $N_j(a, \gamma)$ and, for every $D \geq j$ and $0 < \varepsilon \leq 1$, a positive integer $Q_j(a, \gamma; \varepsilon)$, a number $c_{j,D}(a, \gamma; \varepsilon) \in (0, 1]$, and a threshold $T_{j,D}(a, \gamma; \varepsilon) \geq 2$ with the following property.*

If $p \geq N_j(a, \gamma)$ is prime, $\Gamma \subseteq \mathbb{F}_p^j \times \mathbb{F}_p$ has size at most $\gamma^{-2}p^j$ and has the γ -product property, there is a single $J \subseteq \mathbb{F}_p^j$ with $|J| \geq (1 - a)p^j$ such that $\Gamma \cap (J \times \mathbb{F}_p)$ has the active-coordinate cover of Definition II.235.1 with parameters

$$(Q_j(a, \gamma; \varepsilon), c_{j,D}(a, \gamma; \varepsilon), T_{j,D}(a, \gamma; \varepsilon); \varepsilon)$$

for every D , every permitted active-coordinate map, every ε , and every box whose width meets the threshold. The set J is independent of all these local choices.

The graph count does not require an ambient-dimension index. The width exponent and threshold do. This distinction will follow from the recurrence itself.

For the base case put

$$\begin{aligned} b_* &= 2^{-1882} \gamma^{9312} a^{1165}, & q_* &= \lceil \gamma^{-2}/b_* \rceil, \\ t_* &= \lceil 16q_* b_*^{-2} \rceil. \end{aligned} \tag{II.235.13}$$

For every $D \geq 1$ and every local error, define

$$\begin{aligned} N_1(a, \gamma) &= 3, & Q_1(a, \gamma; \varepsilon) &= q_*, \\ c_{1,D}(a, \gamma; \varepsilon) &= \frac{1}{4(t_* + 1)}, \\ T_{1,D}(a, \gamma; \varepsilon) &= \max \left\{ 2^{4(t_*+1)}, (256\pi/b_*)^{2(t_*+1)} \right\}. \end{aligned} \tag{II.235.14}$$

Proof of the base case. As long as the remaining relation has projection of density at least a , choose one graph over its projection, whose actual density is $v \geq a$. The product property gives at least $\gamma^8 v^4 p^3$ additive quadruples. In I1 its normalized parameter is $\rho = \gamma^8 v$, so the extracted order-eight Freiman piece has size at least $2^{-1882} \gamma^{9312} v^{1165} p \geq b_* p$. Remove the graph of this piece and repeat. There are at most q_* pieces; after deleting the final remaining projection, their union covers the relation above a set J of the required size.

For each piece use its *global* density in I2. The combined spectra have at most t_* frequencies. Each piece's valid Bohr radius is at least the one obtained from b_* . This same finite collection of pieces works in every ambient dimension.

Consider a D -box of width m and common difference d_0 , and put

$$L = \left\lfloor m^{1/(2(t_*+1))} \right\rfloor, \quad H = \lceil 64\pi L/b_* \rceil.$$

Simultaneous pigeonholing for the combined frequencies gives $1 \leq u \leq H^{t*}$ with $|rd_0u|_p \leq p/H$ for every frequency r . The threshold in (II.235.14) implies

$$\frac{m}{u} \geq 2L, \quad L \geq m^{1/(4(t*+1))}.$$

Indeed $H \leq 128\pi L/b_*$, and $m^{1/2} \geq 2(128\pi/b_*)^{t*}$ suffices for the first inequality; the other threshold gives the second after taking the floor.

In each side's index interval split the residue classes modulo u into successive blocks with lengths in $[L, 2L - 1]$. Such a split is possible for every integer length at least L , and every residue class has at least L points. Products of these blocks partition the original box. They have common difference $ud_0 \neq 0$, because $u \leq m/(2L) < p$, and have the required width. The Bohr condition for I2 with length parameter $2L - 1$ follows from

$$\frac{1}{H} \leq \frac{b_*}{64\pi L} \leq \frac{b_*}{32\pi(2L - 1)}.$$

Each graph piece is therefore affine in its active coordinate on every cell. There is no local loss. Translating or choosing a different active coordinate does not change the argument. $\square$

Remark II.235.6 (Why the base proof uses global spectra). The extracted Freiman pieces need not have density b_* in an arbitrary local progression. Thus one cannot simply apply the local-density form of Corollary 7.11 with the original density parameter. The preceding argument first obtains Bohr information from the global sets, then localizes that information by simultaneous pigeonholing.

II.235.6 Global extraction in dimension $k + 1$

Suppose the theorem is known in active dimensions $1, \dots, k$. Fix a, γ . All constants in the following display are global; they are chosen before any ambient dimension or local tolerance:

$$\begin{aligned} \kappa &= a/2^{k+2}, & b_0 &= a/2, \\ \alpha_0 &= b_0^{7 \cdot 4^{k+1} - 15} \gamma^{21k \cdot 4^{k+1}}, \\ \tau &= (2^{-46} \alpha_0)^{H_k} b_0^{15}, & u_0 &= \tau/4, \\ \delta &= 2^{-37} u_0^{11/2}, & R_0 &= \left\lceil 2^{74} u_0^{-10} \right\rceil, \\ \zeta &= 2^{-155 R_0} u_0^{18 R_0} / R_0, \\ \eta &= 2^{-32} \tau^8, & B_* &= \left\lceil \gamma^{-2} / \eta \right\rceil. \end{aligned} \tag{II.235.15}$$

Here H_k was defined in (II.235.5). Keeping the exact Riesz output τ avoids the further coarse numerical domination used in the source's Lemma 15.6. Define the prime threshold by

$$N_{k+1}(a, \gamma) = \left\lceil \max \left\{ 3, \frac{k 3^{2^{k+4}}}{\tau}, \max_{1 \leq l \leq k} N_l(\kappa, \gamma), N_k(\tau/8, \delta) \right\} \right\rceil. \tag{II.235.16}$$

Every dimension on its right side is smaller than $k + 1$.

Lemma II.235.7 (A fixed graph piece). *Assume I1–I7 and the lower-dimensional induction hypotheses. Let p meet (II.235.16), and suppose a γ -product relation of size at most $\gamma^{-2} p^{k+1}$ has projection of density at least a . It contains a graph piece of size at least ηp^{k+1} with the local profiles constructed in Subsection II.235.7. The piece is chosen independently of ambient dimension, local error, and local box.*

Proof. Choose a partial map ϕ on a set A of density at least a , with graph in the relation. For each positive-dimensional proper coordinate section, apply the corresponding lower-dimensional core theorem with outer error κ and parameter γ . For a fixed set of coordinates to be held constant, its sections partition the ambient space, so their total discarded mass is at most κp^{k+1} . There are fewer than 2^{k+1} choices. Pruning all of them leaves a domain of density at least b_0 , on every proper section of which the relevant lower-dimensional core profile holds. Those profiles hold simultaneously in all ambient dimensions and at every local error. They persist on all further restrictions.

Apply I3 and I4 with selection error 2^{-44} . If the actual retained density exceeds b_0 , the same lower bounds remain valid, since all powers of that density are positive. The threshold in (II.235.16) gives a set B_0 containing at least τp^{17k+15} eight-arrangements, at least a proportion $1 - 2^{-44}$ respected by ϕ .

For each side vector h , let $C(h)$ and $G(h)$ count its arrangements and respected arrangements. Let X_h be the set of individual cubes with that side vector, with projection to the last coordinate x . Directly from the ordered arrangement definition,

$$\begin{aligned} C(h) &\leq p^{16k+15}, & |X_h(x)| &\leq p^k, \\ C(h) &\leq |X_h| p^{15k+14}. \end{aligned} \tag{II.235.17}$$

For the last inequality, specify the first cube, the other fifteen basepoints, and fourteen of their last coordinates; the remaining last coordinate is determined by the additive relation.

Discard the h with $G(h) < (1 - 2^{-43})C(h)$. Their total arrangement mass is at most half the total, because the overall failure proportion is at most 2^{-44} . Then discard those with $C(h) < (\tau/4)p^{16k+15}$. The remaining set H satisfies

$$\sum_{h \in H} C(h) \geq (\tau/4)p^{17k+15}, \quad |X_h| \geq u_0 p^{k+1} \quad (h \in H). \tag{II.235.18}$$

The induced cube-sum map on X_h respects at least a proportion $1 - 2^{-43}$ of its order-eight additive relations.

Let $f = \mathbf{1}_{B_0}$, so that the function f_h in I6 is exactly $x \mapsto |X_h(x)|$. Apply I5 with the actual density $b_h = |X_h|/p^{k+1} \geq u_0$. Enlarge its natural spectrum to

$$K_h = \{t : |\hat{f}_h(t)| \geq \delta p^{k+1}\}. \tag{II.235.19}$$

The natural threshold is at least the threshold in this definition. Also $R_0 \geq 2^{74} b_h^{-10}$, and

$$2^{-155R_0} u_0^{18R_0} / R_0 \leq 2^{-155R_{b_h}} b_h^{18R_{b_h}} / R_{b_h}.$$

Consequently the original homomorphism may be restricted to $\text{Bohr}(K_h, \zeta)$. Here R_0 bounds the *natural* spectrum's rank; no bound $|K_h| \leq R_0$ is being asserted for the enlarged spectrum. We obtain $Y_h \subseteq X_h$ with

$$|Y_h| \geq \frac{u_0^6}{20000} |X_h| \geq 2^{-27} \tau^6 |X_h|. \tag{II.235.20}$$

The cube sum on Y_h gives a partial function $\phi'(h, x)$, affine along progressions of length at most L whose difference lies in $\text{Bohr}(K_h, \zeta/L)$.

By I6, the spectral relation $\Delta = \{(h, t) : t \in K_h\}$ has the δ -product property. Parseval and $0 \leq f_h \leq p^k$ imply

$$\sum_t |\hat{f}_h(t)|^2 = p \sum_x f_h(x)^2 \leq p^{2k+2}, \quad |K_h| \leq \delta^{-2}.$$

Thus the lower-dimensional core theorem applies to Δ with outer error $\tau/8$. Intersect its core with H and call the result H_1 . By (II.235.17),

$$\sum_{h \in H_1} C(h) \geq (\tau/8)p^{17k+15}, \quad |H_1| \geq (\tau/8)p^k.$$

Together with (II.235.18) and (II.235.20), this gives

$$\sum_{h \in H_1} |Y_h| \geq 2^{-32} \tau^8 p^{2k+1} = \eta p^{2k+1}. \quad (\text{II.235.21})$$

Average over the p^k possible cube basepoints. There is one fixed $z = (z_1, \dots, z_k)$ and a set B_1 of at least ηp^{k+1} pairs (h, x) , with $h \in H_1$, on which

$$\phi'(h, x) = \sum_{\omega \in \{0,1\}^k} (-1)^{|\omega|} \phi(z + \omega h, x). \quad (\text{II.235.22})$$

All displayed map values are defined there. The all-ones term $\phi_1(h, x) = \phi(z + h, x)$ is a translation of a restriction of ϕ . Solving (II.235.22) for it writes it as $(-1)^k \phi'$ plus $2^k - 1$ signed proper-section terms. Each proper section of ϕ_1 retains its lower-dimensional core profile. All choices made so far concern finite global data. No local box, ambient dimension, or local error has been chosen. The local proof that follows therefore applies to this one fixed piece. $\square$

II.235.7 The exact local profile of an extracted piece

Fix $D \geq k + 1$ and a desired error $0 < e \leq 1$. The local construction has four error allocations, each at most $e/4$.

Spectral freezing. Define

$$\begin{aligned} q_S &= Q_k(\tau/8, \delta; e/4), \\ c_S &= c_{k,D}(\tau/8, \delta; e/4), \quad T_S = T_{k,D}(\tau/8, \delta; e/4), \\ a_S &= a_D(q_S), \quad c_F = c_S a_S / 2, \\ T_F &= \max \left\{ T_S, L_D(q_S)^{1/c_S}, (16/\zeta)^{2/(c_S a_S)} \right\}. \end{aligned} \quad (\text{II.235.23})$$

These expressions use only the k -dimensional spectral core.

Lemma II.235.8 (Freezing in the full ambient box). *On a D -box of width $m \geq T_F$, one may discard at most an $e/4$ fraction and partition into boxes of width at least m^{c_F} such that $\phi'(h, \cdot)$ is affine on every remaining partial vertical row with $h \in H_1$.*

Proof. Apply the active-coordinate spectral cover to the D -box, obtaining cells of width at least m^{c_S} and a good set with loss at most $e/4$. On a cell let its spectral models be $\mu_1(h), \dots, \mu_{q_S}(h)$, padding the list if necessary. They depend only on h . At every good point with $h \in H_1$, the entire set K_h lies among their values.

Apply I7 in ambient dimension D to $\nu_i(h, x) = \mu_i(h)x$. These are multilinear in the ambient coordinates. The resulting boxes have width at least $W = m^{c_S a_S}$ and each ν_i has image diameter at

most $2p/W$. For a child common difference d , compare two consecutive values of its x coordinate, keeping all other coordinates fixed. This gives

$$|\mu_i(h)d|_p \leq 2p/W. \quad (\text{II.235.24})$$

There are two such values because the threshold makes W large.

Set $L = \lfloor \zeta W/8 \rfloor$. Split every side into consecutive blocks of lengths in $[L, 2L - 1]$, leaving the common difference unchanged. Since $W^{1/2} \geq 16/\zeta$,

$$L \geq \zeta W/16 \geq W^{1/2} = m^{c_F}.$$

Also $2/W \leq \zeta/(2L - 1)$. Thus (II.235.24) puts d in the Bohr set required to make $\phi'(h, \cdot)$ affine on a vertical block, whenever the original spectral good set covers K_h there.

The good set is not required to be cylindrical. For any retained point its full K_h is covered by the same cell's models; those models do not depend on x or on unused coordinates. This is enough to establish the direction condition for its row. Refining a good set by intersection with cells creates no further loss. $\square$

The proper-section terms. Write

$$R = 2^k - 1, \quad n_l = \binom{k}{l-1} \quad (1 \leq l \leq k). \quad (\text{II.235.25})$$

Exactly n_l terms in the remainder of (II.235.22) have l active coordinates: $l - 1$ coordinates of h and the distinguished coordinate x . Retaining these dimensions gives better controls than bounding every term by the k -dimensional profile. Put

$$\begin{aligned} q_l &= Q_l(\kappa, \gamma; e/(4R)), \\ c_l &= c_{l,D}(\kappa, \gamma; e/(4R)), \\ T_R &= \max_{1 \leq l \leq k} T_{l,D}(\kappa, \gamma; e/(4R)), \\ q_0 &= \prod_{l=1}^k q_l^{n_l}, \quad c_R = \prod_{l=1}^k c_l^{n_l}. \end{aligned} \quad (\text{II.235.26})$$

Successive refinement of these R functions costs at most $e/4$ of the ambient mass. It changes the width exponent to $c_F c_R$. On each resulting vertical row, ϕ_1 is covered by at most q_0 affine functions: choose one proper-section model for each term and add the single affine function supplied for ϕ' . Assign each point to the first line containing it to get disjoint classes for Lemma II.235.4.

Anchor slices. Define

$$\begin{aligned} r &= \lceil 8q_0/e \rceil, \\ q_H &= Q_k(\kappa, \gamma; e/(4r)), \\ c_H &= c_{k,D}(\kappa, \gamma; e/(4r)), \quad T_H = T_{k,D}(\kappa, \gamma; e/(4r)). \end{aligned} \quad (\text{II.235.27})$$

The anchor lemma loses at most $e/4$ of each ambient cell. For a cell whose vertical side has fewer than r points, take that whole side as anchors. Each anchor slice is a restriction of a proper section of ϕ_1 , so it has the indicated k -dimensional profile *in the same ambient dimension D* , treating the

vertical coordinate as unused. Refine all at most r such slices, for another loss of at most $e/4$. No lifting of an independently chosen lower-dimensional partition is required.

The resulting triple for one extracted piece is

$$\begin{aligned} q_{\text{piece}}(e) &= rq_H + \binom{r}{2} q_H^2, \\ c_{\text{piece}}(e) &= c_F c_R c_H^r, \\ T_{\text{piece}}(e) &= \max\{T_F, T_R^{1/c_{\text{piece}}(e)}, T_H^{1/c_{\text{piece}}(e)}\}. \end{aligned} \tag{II.235.28}$$

The threshold is deliberately generous: every input width in the remainder and anchor refinements is at least $m^{c_{\text{piece}}(e)}$. The spectral threshold was already accounted for in T_F . The four losses sum to at most e . Formula (II.235.12) shows that all final models are multilinear in the active $k+1$ coordinates. This completes the local assertion of Lemma II.235.7.

II.235.8 Closing the induction

For $j = k+1$, use the global constants (II.235.15). Given ε , set $e = \varepsilon/B_*$ and define

$$\begin{aligned} Q_{k+1}(a, \gamma; \varepsilon) &= B_* q_{\text{piece}}(e), \\ c_{k+1,D}(a, \gamma; \varepsilon) &= c_{\text{piece}}(e)^{B_*}, \\ T_{k+1,D}(a, \gamma; \varepsilon) &= \max\{2, T_{\text{piece}}(e)^{1/c_{k+1,D}(a, \gamma; \varepsilon)}\}. \end{aligned} \tag{II.235.29}$$

Equations (II.235.13)–(II.235.14), (II.235.15)–(II.235.16), and (II.235.23)–(II.235.29) are the promised explicit recurrences.

Completion of Theorem II.235.5. Apply Lemma II.235.7 whenever the remaining relation has projection of density at least a , and remove the resulting graph piece. Each removes at least ηp^{k+1} graph points. The total number of pieces is at most B_* , since the original relation has at most $\gamma^{-2} p^{k+1}$ points. Removing a graph preserves the product property. At termination let J be the complement of the remaining projection. Then $|J| \geq (1-a)p^{k+1}$, and the relation above J is covered by the union of the extracted pieces.

The lower-dimensional core choices, Riesz selections, domain subsets, and cube basepoints all precede the choice of a local parameter. Thus these pieces and J work simultaneously for every ambient dimension and local tolerance. Now fix those local choices. Use each piece at error ε/B_* , refine successively, and apply Lemma II.235.3. The graph counts add, the width exponents multiply, and the errors sum to at most ε . Padding by empty pieces if necessary gives exactly the bounds in (II.235.29).

Every profile called on the right side has active dimension at most k . There is no same-dimension call whose size must be absorbed into a previously fixed power. The recurrences are therefore well founded. At the base Q_1 does not depend on D ; the graph-count recurrence uses only lower graph counts, whereas the spectral freezing stage contributes one affine map per row independently of its number of spectral models. Hence Q_{k+1} also does not depend on D . This finishes the induction. $\square$

II.235.9 Growth bounds for the new controls

Proposition II.235.9 (Elementary envelopes). *For every fixed j, D , $D \geq j$, there is an effectively computable $C \geq 2$ such that, with $u = a\gamma\varepsilon$,*

$$\begin{aligned} N_j(a, \gamma) &\leq (C/(a\gamma))^C, \\ Q_j(a, \gamma; \varepsilon) &\leq (C/u)^C, \\ c_{j,D}(a, \gamma; \varepsilon) &\geq \exp(-(C/u)^C), \\ T_{j,D}(a, \gamma; \varepsilon) &\leq \exp\left(\exp((C/u)^C)\right). \end{aligned} \tag{II.235.30}$$

The exact recurrences, rather than a claim of an optimal C , are the numerical content of the theorem.

Proof. Use induction on the active dimension. The base constants have polynomial upper bounds in $(a\gamma)^{-1}$; their negative width logarithm has such a bound as well. At an inductive step, the quantities $\kappa^{-1}, \tau^{-1}, \delta^{-1}, \eta^{-1}, B_*$ are polynomially bounded in $(a\gamma)^{-1}$ for fixed dimension. So is $\log(1/\zeta)$: R_0 is polynomially bounded and the only additional factors in its logarithm are logarithms of such quantities. The recursion for N_j therefore remains polynomial.

Every lower-dimensional parameter substituted into a Q is at least the reciprocal of a polynomial in u^{-1} . Consequently the q_l, q_0, r, q_H , and the final graph count are polynomially bounded. The negative logarithm of a width exponent is a sum of lower negative logarithms multiplied by polynomially bounded repetition counts. In the spectral step, $-\log a_D(q_S) = 2^D q_S \log K_D$ is also polynomially bounded. This proves the asserted single-exponential lower bound for c .

Finally, the logarithm of each threshold is bounded by an exponential of a polynomial. This is true of $L_D(q_S)$ and of the lower thresholds by induction. Raising such a threshold to the reciprocal of a width exponent multiplies its logarithm by another single exponential, preserving this form. The term involving $1/\zeta$ obeys the same bound. Increasing C finitely many times proves all four inequalities simultaneously. This induction can be performed directly on the explicit expressions, so the constant is effective. $\square$

The effect of preserving actual section dimensions is visible even when the outer parameters are fixed. Define integers A_j, B_j by

$$\begin{aligned} A_1 &= B_1 = 0, & S_k &= \sum_{l=1}^k \binom{k}{l-1} A_l, \\ A_{k+1} &= 2(A_k + 1)(S_k + 2) - 2, \\ B_{k+1} &= (B_k + 1)(S_k + 2) - 1. \end{aligned} \tag{II.235.31}$$

Proposition II.235.10 (Dependence on the local error). *For fixed a, γ, j, D , as $\varepsilon \downarrow 0$ the exact recurrences satisfy*

$$Q_j = O(\varepsilon^{-A_j}), \quad -\log c_{j,D} = O(\varepsilon^{-B_j}), \quad \log \log T_{j,D} = O(\varepsilon^{-B_j}). \tag{II.235.32}$$

The first four pairs (A_j, B_j) are $(0, 0), (2, 1), (34, 11), (7698, 1319)$.

Proof. The statement is immediate in dimension one. At the next step, $q_0 = O(\varepsilon^{-S_k})$, so $r = O(\varepsilon^{-(S_k+1)})$, and the horizontal error $e/(4r)$ is bounded below by a positive constant times ε^{S_k+2} .

The graph bound $rq_H + \binom{r}{2}q_H^2$ has exponent $2(S_k + 1) + 2A_k(S_k + 2) = A_{k+1}$. The term c_H^r has negative logarithm of order at most

$$\varepsilon^{-(S_k+1)-B_k(S_k+2)} = \varepsilon^{-B_{k+1}}.$$

All earlier width losses have exponents at most $\max\{A_k, B_k\} \leq B_{k+1}$. For the thresholds, the lower horizontal threshold contributes exponent $B_k(S_k + 2)$ to its double logarithm, and taking reciprocal width powers contributes at most B_{k+1} , which is larger. The other thresholds are smaller in the same sense. The fixed number B_* of global pieces changes constants but none of these powers. $\square$

Replacing all proper-section terms by the largest-dimensional profile would instead give the pairs $(0, 0), (2, 1), (46, 15), (30454, 5183)$. The smaller exponents above are proved for this recurrence; no optimality is asserted.

II.235.10 A replacement for the multilinear inverse conclusion

Corollary II.235.11 (Recursive version of the Corollary 16.11 interface). *Assume I1–I7 and the large-coefficient consequence of the paper’s Lemma 3.1. Let $0 < \alpha \leq 1/2$ and $f : \mathbb{F}_p \rightarrow \mathbb{C}$, $|f| \leq 1$, fail α -uniformity of degree $k + 1$ in the paper’s convention. Put*

$$a = \alpha/4, \quad \gamma = \alpha/2, \quad \varepsilon = \alpha/8,$$

and abbreviate $Q = Q_k(a, \gamma; \varepsilon)$ and $c = c_{k,k}(a, \gamma; \varepsilon)$. If

$$p \geq \max\{N_k(a, \gamma), T_{k,k}(a, \gamma; \varepsilon)\}, \tag{II.235.33}$$

there is a box $P \subseteq \mathbb{F}_p^k$ of width at least p^c and a multilinear polynomial μ on P such that

$$\#\{h \in P : |\widehat{\Delta_h f}(\mu(h))| \geq (\alpha/2)p\} \geq \frac{\alpha}{8Q}|P|. \tag{II.235.34}$$

Here $\Delta_h f$ is the k -fold multiplicative derivative.

Proof. The large-coefficient step gives a partial graph on at least $(\alpha/2)p^k$ side vectors, selecting Fourier coefficients of modulus at least $(\alpha/2)p$. By I6 this graph has the $(\alpha/2)$ -product property. Apply Theorem II.235.5 and then its local conclusion to the whole box $\mathbb{F}_p^k$, of width p . Global deletion removes at most $(\alpha/4)p^k$ points of this graph’s domain; local deletion removes at most $(\alpha/8)p^k$. At least $(\alpha/8)p^k$ domain points therefore survive. Some partition cell has surviving relative density at least $\alpha/8$. Assign each of its surviving points to the first model whose value is the selected coefficient. One of at most Q models receives at least the fraction $\alpha/(8Q)$, proving (II.235.34). $\square$

[write] **7 October 2026; sources 47–53.** *Against the formal statements* (checked at HEAD 06b30ca52). `corollary_16_11` (Section16.lean, line 893) and `lemma_16_10` (line 851) are open; the repository has proved `corollary_16_11` conditionally on its Section 16 induction (`corollary_16_11_of_theorem_16_2`, `Proofs16CorollaryFromInduction.lean`, line 123) and reduced Theorem 16.2 to unasserted contextual lifting statements (`Section16ContextualLiftAt`, `Proofs16ContextualInduction.lean`, line 18) and, since c92c64d18, to an unasserted selected-lift obligation (`Section16SelectedLiftAt`, `Proofs16SelectedLiftIn` `duction.lean`, line 14; `theorem_16_2_of_selected_lifting`, line 64). This corollary is a different, written route with its own controls, conditional on I1–I7; it is named in `FORMALIZATION_STATUS.txt` as a candidate for the repaired Lemma 16.10, with source 50’s interface. Nothing in this section is formalized.

[write] **7 October 2026; sources 60–69.** At *HEAD fb9602e55* (checked at the eleventh write): `Section16ContextualLiftAt` is now at line 20 of `Proofs16ContextualInduction.lean` and `theorem_16_2_of_selected_lifting` at line 65 of `Proofs16SelectedLiftInduction.lean`; the universal contextual-lift contract has been proved false (`section16_contextual_lift_counterexample`, `Proofs16ContextualLiftCounterexample.lean`, line 80; 48efbb992), following source 87, which is placed and not yet written in. `theorem_16_2` and `lemma_16_10` remain open; the existential selected-lifting obligation is unchanged.

If a common width and density parameter is wanted, one may take $c_* = \min\{c, \alpha/(8Q)\}$. More informatively, Proposition II.235.9 gives a polynomial density lower bound in α and a width of the form

$$p^{\exp(-(C/\alpha)^C)}$$

once $p \geq \exp(\exp((C/\alpha)^C))$, with effective C depending on k . This is a complete recursive replacement of the local implication, with a weaker width exponent than the unsupported printed power. Every use of a source interface and every new quantitative dependency is displayed above. A further claim about the downstream Sections 17–18 would require propagating these new controls through those sections; that is not part of this theorem.

II.235.11 Questions arising from the repaired recurrence

The proof suggests several focused problems. First, can the repeated slice refinement be replaced by a simultaneous theorem that keeps the width exponent polynomial in the local error? The present recurrence identifies precisely where the single exponential appears: a number of anchor slices growing as a power of ε^{-1} is refined successively. Second, can the minimum-width threshold be reduced while preserving the global core’s independence of the local tolerance? Third, one can replace I7 by a sharper simultaneous polynomial recurrence theorem and recompute the profiles without modifying the global extraction argument. Finally, a formal implementation can represent the recurrences as expression trees indexed by active dimension; the strict decrease in every recursive call gives a direct termination proof and avoids encoding a failed final domination as an opaque numerical assertion.

II.236 Sources 55 and 57: the polynomial-phase endpoint in every degree [write]

Added in the tenth write, 7 October 2026.

The next eleven sections are source 55’s Sections 2–11 and source 57’s Section 5, on functions whose Gowers norm is close to its maximum. *Logically they continue the chapters on the Proposition 6.1 endpoint* of sources 31 and 29 (Section II.104) and 38 (Section II.149); they are printed here because no number printed earlier may move. Source 55’s introduction is Section F.60; source 57’s is Section F.62, and its Section 13 chapter is Section II.287.

What source 55 proves. For $|f| \leq 1$ on a finite abelian group G and $\varepsilon = 1 - \|f\|_{U^d}^{2^d}$, the squared normalized L^2 distance to the phase polynomials of degree at most $d - 1$ (nonclassical ones included) satisfies $D_d(f) \leq (2/q)\varepsilon + b_d(G)\varepsilon^2 + R_d\varepsilon^3$ once $\varepsilon \leq e_d$, $q = 2^d$, with $b_d = (3^d - 1)/(2q^2)$ for $|G|$ odd and $(q^2 - 1)/(3q^2)$ for $|G|$ even, both sharp on every nontrivial group of that parity (Theorem II.237.1); second-order saturation forces the tangent to be a conjugate pair of characters (odd order) or one order-two character (even order) (Theorem II.237.3). The local theorem is explicit (radius (II.242.5), cubic constant (II.242.6); Theorem II.242.1); the global entry e_d comes from the

published Eisner–Tao theorem (Theorem II.243.1), which source 55 imports and does not make numerical. The key steps are a fifth-order real-part cancellation for bounded complex functions and a four-vertex census of the cube (Lemma II.240.2, Lemma II.239.2); the exact even-order U^2 modulus is $2(1 - \sqrt{(1 + \sqrt{1 - 2e})/2})$ (Proposition II.245.1). At $d = 3$ the coefficients 13/64 and 21/64 are those of sources 38 and 31 in another class of functions (note after Theorem II.237.1).

What source 57's Section 5 proves. If $Q_4(f) \geq 1 - \varepsilon$ with $\varepsilon \leq 2^{-10}$ on $\mathbb{F}_p$, one scalar c gives $|\widehat{\Delta_h \Delta_k f}(chk)| \geq 3/4 - 4\varepsilon$ on at least $(1 - 332\varepsilon)p^2$ pairs (Theorem II.247.5), with no scale threshold and no inverse theorem as input; at every order $216s + 8$ in place of 332 (Theorem II.247.7), an s -additive frequency map on every finite abelian group (Theorem II.247.9), a cubic phase with correlation at least $\sqrt{1 - 6194\varepsilon} - \varepsilon$ in characteristic at least five (Corollary II.247.10), polynomial phases of degree $k - 1$ for $p \geq k$ (Corollary II.247.11), and the linear order of the squared error is optimal (Proposition II.247.12). This is the endpoint branch of its proof of Theorem 13.12 (Section II.287).

Relation to the formal statements (checked at HEAD fe7165a3c). `proposition_6_1` (Sections06_07.lean, line 37) has the exact companion `proposition_6_1_holds` (Proofs05_10.lean, line 505); no declaration concerns near-maximal Gowers norms in general. Source 57's Theorem II.247.5 is formalized as the repository's own `endpoint_global_fourier` (Proofs13EndpointGlobalFourier.lean, line 12; density $1 - 332\varepsilon$, coefficient $(3/4 - 4\varepsilon)N$) for the proof of `theorem_13_12` (note in Section II.247); nothing else in this chapter is formalized.

II.237 Main results

Source 55, delivered Section 2.

Theorem II.237.1 (Sharp second-order endpoint stability). *For every integer $d \geq 2$ there are constants $e_d > 0$ and $R_d < \infty$, independent of the finite abelian group, such that for every $|f| \leq 1$ with $\varepsilon = 1 - \mathcal{U}_d(f) \leq e_d$,*

$$D_d(f) \leq \frac{2}{q}\varepsilon + b_d(G)\varepsilon^2 + R_d\varepsilon^3, \quad (\text{II.237.1})$$

where

$$b_d(G) = \begin{cases} \frac{3^d - 1}{2q^2}, & |G| \text{ odd}, \\ \frac{q^2 - 1}{3q^2}, & |G| \text{ even}. \end{cases} \quad (\text{II.237.2})$$

For every nontrivial finite abelian G ,

$$\mathfrak{S}_{d,G}(e) = \frac{2}{q}e + b_d(G)e^2 + O_d(e^3) \quad (e \downarrow 0), \quad (\text{II.237.3})$$

with a range and error constant uniform in G . In particular, neither the linear nor the quadratic coefficient can be decreased.

[write] **7 October 2026; sources 54–59.** *Checked at the write* (exact rationals). With $C_d = (6^d - 2q^2 + q)/8$, $A_d = q(q-1)(q-2)/24$ and $K_{0,d} = q(q-1)/8$ (Appendix V.210), $K_{0,d} + C_d/2 = q(3^d - 1)/16$ and $K_{0,d} + A_d = q(q^2 - 1)/24$ for $d = 2, \dots, 12$, so $8K_d/q^3$ is (II.237.2); the table above and $\beta_d(G)$ of Corollary II.237.2 follow for $d = 2, \dots, 6$. *Against sources 31 and 38.* Source 31's optimal fidelity bound $R(\delta) = (1 - \sqrt{(1 + \sqrt{1 - 4\delta})/2})/2$ gives the squared distance $2(1 - \sqrt{1 - R(\delta)}) = \delta/4 + 21\delta^2/64 + O(\delta^3)$ (Theorem II.107.4, uniform over finite

abelian groups), and source 38's odd-order modulus has squared distance $\varepsilon/4 + 13\varepsilon^2/64 + (\sqrt{2}/64)\varepsilon^{5/2} + O(\varepsilon^3)$ (Theorem II.150.11); at $d = 3$ this theorem's coefficients are $1/4$, $13/64$ (odd order) and $21/64$ (even order). The classes differ: sources 31 and 38 take L^2 -normalized functions and coset phases, this source $|f| \leq 1$ and phase polynomials on the whole group. Unimodular functions belong to both, and this source's extremizing families ((II.243.2), (II.243.6)) are unimodular; source 38's $\varepsilon^{5/2}$ term comes from near point masses on groups with three-torsion, which are not bounded by one after normalization, so it does not contradict the $O_d(\varepsilon^3)$ here. This source extends the two coefficients to every degree and to bounded functions; the global entry e_d rests on the published Eisner–Tao theorem and is not explicit (item II.2).

An explicit admissible R_d is given in equation (II.242.6); the local radius is explicit in equation (II.242.5). We do not extract a numerical value of the global entry threshold e_d from the published near-extremizer theorem. That distinction does not make theorem II.237.1 conditional: its input is a published theorem, not an unproved hypothesis.

Corollary II.237.2 (Norm-deficit normalization). *Writing $\delta = 1 - \|f\|_{U^d}$ instead, the sharp expansion is*

$$\mathfrak{S}_{d,G}(1 - (1 - \delta)^q) = 2\delta + \beta_d(G)\delta^2 + O_d(\delta^3),$$

where

$$\beta_d(G) = \begin{cases} (3^d - 2^{d+1} + 1)/2, & |G| \text{ odd}, \\ (2^d - 1)(2^d - 2)/3, & |G| \text{ even}. \end{cases}$$

Proof. Substitute $1 - (1 - \delta)^q = q\delta - q(q - 1)\delta^2/2 + O_d(\delta^3)$ into theorem II.237.1 and its matching examples. $\square$

For orientation, the first coefficients in the cubical-defect normalization are as follows.

d	coefficient of e	odd-order coefficient of e^2	even-order coefficient of e^2
2	$1/2$	$1/4$	$5/16$
3	$1/4$	$13/64$	$21/64$
4	$1/8$	$5/32$	$85/256$
5	$1/16$	$121/1024$	$341/1024$
6	$1/32$	$91/1024$	$1365/4096$

Theorem II.237.3 (Tangent rigidity of second-order extremizers). *Fix $d \geq 2$. Let G_n be nontrivial finite abelian groups of the same parity, and let $|f_n| \leq 1$, with $\varepsilon_n = 1 - \mathcal{U}_d(f_n) > 0$ and $\varepsilon_n \rightarrow 0$. Suppose*

$$D_d(f_n) = \frac{2}{q}\varepsilon_n + b_d(G_n)\varepsilon_n^2 + o(\varepsilon_n^2). \quad (\text{II.237.4})$$

Choose minimizing phases P_n , rotated so that $\mathbb{E}f_n\overline{P_n} \geq 0$. Then

$$1 - \mathbb{E}|f_n|^2 = o(\varepsilon_n^2).$$

If the groups have even order, there are nontrivial order-two characters χ_n and signs $s_n \in \{-1, 1\}$ such that

$$\left\| \frac{f_n\overline{P_n} - 1}{\sqrt{D_d(f_n)}} - is_n\chi_n \right\|_2 \longrightarrow 0. \quad (\text{II.237.5})$$

If the groups have odd order, there are nontrivial characters χ_n and $\zeta_n \in S^1$ such that

$$\left\| \frac{f_n\overline{P_n} - 1}{\sqrt{D_d(f_n)}} - i\sqrt{2} \operatorname{Re}(\zeta_n\chi_n) \right\|_2 \longrightarrow 0. \quad (\text{II.237.6})$$

This is a necessity statement for second-order saturation, not a claim that an arbitrary $o(\sqrt{D})$ approximation to the displayed tangent automatically saturates the quadratic coefficient. Quantitative control of higher-order components is needed for that converse. Explicit saturating families will be given in section II.243.

II.238 Demodulation and the bounded-function geometry

Source 55, delivered Section 3.

Fix a phase P , not necessarily minimizing, and rotate it so that $F = f\bar{P}$ has real nonnegative mean. For this section write

$$m = \mathbb{E}F, \quad r = F - m, \quad a = 1 - \mathbb{E}|F|^2, \quad D = \mathbb{E}|F - 1|^2, \quad z = D + a. \quad (\text{II.238.1})$$

The letter r denotes a centered function, not a cube direction. We will use

$$m = 1 - \frac{z}{2}, \quad \mathbb{E}r = 0, \quad v := \mathbb{E}|r|^2 = D - \frac{z^2}{4}, \quad 0 \leq v \leq D. \quad (\text{II.238.2})$$

The identities follow by expanding the squared distances and using $\mathbb{E}|F|^2 = 1 - a$. The number a is an amplitude deficit; it vanishes precisely when $|F| = 1$ everywhere.

Lemma II.238.1 (Amplitude loss is visible to the cube). *For every $d \geq 1$ and every $|F| \leq 1$,*

$$0 \leq a \leq 1 - \mathcal{U}_d(F).$$

Proof. Any two distinct cube vertices are independent and uniform on G . Taking absolute values in (F.60.1) and discarding all but two factors gives

$$\mathcal{U}_d(F) \leq (\mathbb{E}|F|)^2 \leq \mathbb{E}|F|^2.$$

□

The next pointwise identity will rule out an otherwise possible error of order $D^{5/2}$.

Lemma II.238.2 (Real-part constraint). *Let $A(x) = 1 - |F(x)|^2 \geq 0$ and $u = \text{Re } r$. If $m > 0$, then*

$$2mu = v + a - A - |r|^2. \quad (\text{II.238.3})$$

For every nontrivial character χ ,

$$\left| \widehat{r}(\chi) + \overline{\widehat{r}(-\chi)} \right| \leq z. \quad (\text{II.238.4})$$

Here $\widehat{r}(\chi) = \mathbb{E}r\bar{\chi}$, and additive notation is used in the character group.

Proof. Expand $|F|^2 = |m + r|^2$ and use $1 - m^2 = v + a$ to obtain (II.238.3). Also

$$2(1 - \text{Re } F) = A + |F - 1|^2.$$

The right side is nonnegative and has mean $a + D = z$. At a nontrivial frequency its Fourier coefficient has magnitude at most z . Since constants disappear at that frequency and $2\text{Re } \widehat{r}(\chi) = \widehat{r}(\chi) + \overline{\widehat{r}(-\chi)}$, this proves (II.238.4). □

In particular, a small bounded perturbation of 1 is primarily imaginary. Crucially, (II.238.3) is an exact pointwise identity, so it remains useful for perturbations concentrated on very small subsets. No uniform pointwise smallness of $F - 1$ will be assumed.

II.239 Four vertices of a discrete cube

Source 55, delivered Section 4.

For $\omega \in \{0, 1\}^d$, put $X_\omega = x + \omega \cdot h$, with all x, h_j independent and uniform on G .

II.239.1 Independence of every three vertices

Lemma II.239.1 (Three-wise independence). *For any three distinct cube vertices $\omega_1, \omega_2, \omega_3$, the variables $X_{\omega_1}, X_{\omega_2}, X_{\omega_3}$ are independent and uniform, on every finite abelian group.*

Proof. Reflect coordinates to send ω_1 to 0. This is realized on the sampling variables by replacing suitable h_j with $-h_j$ and translating x , an invertible transformation preserving the uniform measure. The remaining two vertices are distinct nonzero binary vectors u, v . Their coordinate columns (u_j, v_j) contain at least two distinct nonzero members of $\{(1, 0), (0, 1), (1, 1)\}$. The corresponding 2×2 determinant is 1 or -1 . Including the base variable x therefore gives a 3×3 minor of determinant ± 1 in the augmented forms $(1, \omega_i)$. Fixing every other variable, these three variables can be solved uniquely for any prescribed triple of outputs over G . $\square$

II.239.2 The complete four-vertex classification

A four-set of cube vertices will be called *ordinary* if it is an integer parallelogram, that is, if its four distinct elements can be ordered so that $a + b = c + e$ in $\mathbb{Z}^d$. It will be called *exceptional* if its xor is zero but it is not ordinary. These are precisely the additional four-point dependencies created by two-torsion.

Lemma II.239.2 (Integral and binary four-point circuits). *Every four-set of distinct vertices belongs to exactly one of the following classes.*

- (i) *An ordinary parallelogram. Three associated variables are independent; the fourth satisfies their unique parallelogram relation.*
- (ii) *An exceptional set. After reflecting and relabeling, its joint law is*

$$(x, x + U + V, x + U + W, x + V + W), \quad (\text{II.239.1})$$

where x, U, V, W are independent uniform elements of G .

- (iii) *A set with nonzero xor. All four associated variables are independent and uniform on G .*

The numbers of ordinary, total xor-zero, and exceptional four-sets are respectively

$$C_d = \frac{6^d - 2 \cdot 4^d + 2^d}{8}, \quad (\text{II.239.2})$$

$$A_d = \frac{q(q-1)(q-2)}{24}, \quad (\text{II.239.3})$$

$$J_d = A_d - C_d. \quad (\text{II.239.4})$$

Proof. Reflect one vertex to 0, and call the others u, v, w . If $u + v + w \neq 0$ in $\mathbb{F}_2^d$, these three distinct nonzero vectors are linearly independent over $\mathbb{F}_2$. There is an odd 3×3 binary minor. A

3×3 zero-one determinant has absolute value at most 2, so this minor is ± 1 . The augmented four forms consequently have a unimodular minor, which proves (iii). For completeness, if B is a 3×3 zero-one matrix, form a 4×4 matrix whose first row and column are all 1 and whose lower block is $1 - 2B$. Subtracting its first row from the other rows gives determinant $(-2)^3 \det B$. Each row has Euclidean norm 2, so the Hadamard determinant inequality gives $8|\det B| \leq 16$.

If $u + v + w = 0$ in $\mathbb{F}_2^d$, every nonzero coordinate column is one of 110, 101, 011. At least two of these types occur. If exactly two occur, one of u, v, w is the integer sum of the other two, whose supports are disjoint. This is an ordinary parallelogram. The independence assertion follows from lemma II.239.1. If all three types occur, combine all sampling directions of each type into a single variable. The three sums are independent and uniform, and give (II.239.1). The nonzero full minors then have greatest common divisor 2; this also describes the relevant Smith obstruction without requiring the Smith normal form theorem in the proof.

There are 6^d ordered solutions of $a + b = c + e$ in the binary cube: each coordinate has six possibilities. The two diagonal families have size 4^d each and intersection of size 2^d . All other solutions have four distinct entries. Each ordinary four-set contributes eight orderings, proving (II.239.2).

Finally, a xor-zero four-set is an affine two-plane in $\mathbb{F}_2^d$. Any ordered triple of distinct points determines its fourth point, and each plane has $4 \cdot 3 \cdot 2 = 24$ such triples. This proves (II.239.3) and (II.239.4). $\square$

d	$\binom{2^d}{4}$	ordinary C_d	exceptional J_d	independent
2	1	1	0	0
3	70	12	2	56
4	1820	100	40	1680
5	35960	720	520	34720

The ordinary/exceptional distinction concerns the integer coefficient matrix. Even in an odd-order group, the exceptional four-set exists combinatorially; its induced map is then surjective because multiplication by 2 is invertible. Its centered correlation consequently vanishes.

II.239.3 Fourier control of the quartic coefficient

For a centered complex function r , define

$$Q_d(r) = \sum_{\substack{S \subseteq \{0,1\}^d \\ |S|=4}} \mathbb{E} \prod_{\omega \in S} c^{|\omega|} r(X_\omega), \quad (\text{II.239.5})$$

$$E_4(r) = \sum_{\chi \in \widehat{G}} |\widehat{r}(\chi)|^4, \quad T_4(r) = \sum_{\substack{\chi \neq 0 \\ 2\chi=0}} |\widehat{r}(\chi)|^4. \quad (\text{II.239.6})$$

Proposition II.239.3 (Quartic circuit bound). *For every centered complex r ,*

$$\operatorname{Re} Q_d(r) \leq C_d E_4(r) + J_d T_4(r). \quad (\text{II.239.7})$$

If b is real and centered, then

$$Q_d(b) = Q_d(ib) = C_d E_4(b) + J_d T_4(b). \quad (\text{II.239.8})$$

Proof. Independent four-sets contribute zero. For an ordinary set, Fourier expansion and the parallelogram relation leave one freely varying frequency, with signs and conjugations depending on the labeling. The resulting sum is a product of four Fourier coefficients, each of the form $\hat{r}(\chi)$, $\hat{r}(-\chi)$, or its conjugate. Hölder's inequality bounds its absolute value by $E_4(r)$.

For (II.239.1), Fourier orthogonality forces

$$\chi_1 + \chi_2 = \chi_1 + \chi_3 = \chi_2 + \chi_3 = 0, \quad \chi_0 + \chi_1 + \chi_2 + \chi_3 = 0.$$

Thus all four frequencies equal the same χ with $2\chi = 0$. The zero frequency contributes nothing because r is centered. The absolute contribution is at most $T_4(r)$. Summation proves (II.239.7).

For real b , an ordinary correlation is exactly $E_4(b)$. On a two-torsion character the Fourier coefficient of a real function is real, so the exceptional correlation is exactly $T_4(b)$. Conjugations do not change b . Replacing b by ib multiplies a four-set contribution by $i^4(-1)^{\sum_{\omega \in S} |\omega|}$. Every dependent four-set has xor zero, so this sign is 1. Independent contributions remain zero. $\square$

II.240 A cubic remainder for bounded complex perturbations

Source 55, delivered Section 5.

An ordinary application of Hölder to five small L^2 factors suggests an error of order $v^{5/2}$, which is too large for a cubic error in the defect. The pointwise geometry of the unit disk improves that estimate.

II.240.1 Hölder estimates from triple independence

Lemma II.240.1 (Four-, five-, and six-factor estimates). *Let $r : G \rightarrow \mathbb{C}$, put $v = \mathbb{E}|r|^2$, and use distinct cube vertices for the variables displayed below. Then*

$$\mathbb{E} \prod_{j=1}^4 |r(X_j)| \leq v^2, \tag{II.240.1}$$

$$\mathbb{E}|r(X_1)|^2 \prod_{j=2}^5 |r(X_j)| \leq v^3, \tag{II.240.2}$$

$$\mathbb{E} \prod_{j=1}^6 |r(X_j)| \leq v^3. \tag{II.240.3}$$

For every nonnegative A with $\mathbb{E}A = a$,

$$\mathbb{E}A(X_1) \prod_{j=2}^5 |r(X_j)| \leq av^2. \tag{II.240.4}$$

Statements with more vertices than the cube contains are simply unused.

Proof. Pair the four factors in (II.240.1) and use Cauchy–Schwarz and pairwise independence. For (II.240.2), take the six triples $\{1, j, k\}$ with $2 \leq j < k \leq 5$. The product of

$$(|r(X_1)|^2 |r(X_j)|^2 |r(X_k)|^2)^{1/6}$$

over these triples is precisely the integrand on the left of (II.240.2). Hölder with six equal exponents and lemma II.239.1 gives v^3 . Replace $|r(X_1)|^2$ by $A(X_1)$ in these six factors to get (II.240.4).

For (II.240.3), use all twenty triples from six indices, with exponent $1/20$ on each product of three squared magnitudes. Each index occurs ten times, giving exponent one on its magnitude. Triple independence again gives the asserted bound. $\square$

II.240.2 Cancellation of the potentially dangerous fifth term

Lemma II.240.2 (Fifth-order real-part cancellation). *In the notation (II.238.1)–(II.238.2), suppose $m > 0$. For any five distinct cube vertices and any conjugation choices $s_j \in \{0, 1\}$,*

$$\left| \operatorname{Re} \mathbb{E} \prod_{j=1}^5 C^{s_j} r(X_j) \right| \leq \frac{5}{m} (v + a) v^2. \quad (\text{II.240.5})$$

Proof. Write $r = u + ib$ with real u, b . A product of five factors each equal to ib or $-ib$ is purely imaginary. Telescope the difference from that product, changing one factor at a time. Since $|b| \leq |r|$, its real part has absolute value at most

$$\sum_{j=1}^5 \mathbb{E} |u(X_j)| \prod_{k \neq j} |r(X_k)|.$$

By (II.238.3),

$$|u| \leq \frac{v + a + A + |r|^2}{2m}.$$

For each summand, lemma II.240.1 bounds the three contributions in the numerator by $(v + a)v^2$, av^2 , and v^3 . Their sum is $2(v + a)v^2$. Summing over the five choices proves (II.240.5). $\square$

Theorem II.240.3 (Explicit centered expansion). *Suppose $|F| \leq 1$ and $m = \mathbb{E}F$ is real with $m \geq 1/2$. In the notation above,*

$$\mathcal{U}_d(F) = m^q + m^{q-4} \operatorname{Re} Q_d(r) + \mathcal{R}, \quad |\mathcal{R}| \leq B_d (v + a) v^2, \quad (\text{II.240.6})$$

where

$$B_d = 10 \binom{q}{5} + \sum_{s=6}^q \binom{q}{s} 2^{s-6}. \quad (\text{II.240.7})$$

A binomial coefficient with lower index larger than q is zero, and an empty sum is zero. In particular, $B_2 = 0$ and $B_3 = 608$.

Proof. Expand each cube factor as $m + \mathcal{C}^{|\omega|} r(X_\omega)$. Terms with one, two, or three centered factors vanish by lemma II.239.1. The constant term is m^q and the four-factor term is $m^{q-4} Q_d(r)$. Taking real parts is legitimate because $\mathcal{U}_d(F)$ is real.

For each five-factor term, lemma II.240.2 and $m \geq 1/2$ give the bound $10(v + a)v^2$; the outside factor m^{q-5} is at most one whenever such terms occur. For a term with $s \geq 6$ factors, retain any six and bound the others using $|r| \leq |F| + m \leq 2$. Equation (II.240.3) gives $2^{s-6} v^3 \leq 2^{s-6} (v + a) v^2$. Summing the numbers of terms yields (II.240.7). For $q = 4$ there is no remainder. $\square$

The estimate is uniform over groups and does not invoke an inverse theorem. It applies to sparse perturbations, amplitude defects, and mixtures of the two. Its constants are deliberately coarse; the point is the order of the error and its explicit dependence only on d .

II.241 A continuous two-torsion profile

Source 55, delivered Section 6.

Let

$$v_2 = \sum_{\substack{\chi \neq 0 \\ 2\chi = 0}} |\hat{r}(\chi)|^2.$$

This is the part of the centered variance carried by nontrivial order-two characters. In odd order $v_2 = 0$.

Proposition II.241.1 (Fourier mass pairing). *Under the bounded-function hypotheses of section II.238,*

$$E_4(r) \leq v_2^2 + \frac{1}{2}(v - v_2)^2 + z^2(v - v_2), \quad (\text{II.241.1})$$

$$T_4(r) \leq v_2^2. \quad (\text{II.241.2})$$

Consequently,

$$\operatorname{Re} Q_d(r) \leq A_d v_2^2 + \frac{C_d}{2}(v - v_2)^2 + C_d z^2(v - v_2). \quad (\text{II.241.3})$$

Also, without using the real-part constraint,

$$\operatorname{Re} Q_d(r) \leq A_d v^2. \quad (\text{II.241.4})$$

Proof. Partition the nonzero characters not fixed by negation into pairs $\{\chi, -\chi\}$. For one pair write $x = |\hat{r}(\chi)|$, $y = |\hat{r}(-\chi)|$, and $s = x^2 + y^2$. Equation (II.238.4) gives $|x - y| \leq z$, whence

$$x^4 + y^4 = \frac{s^2}{2} + \frac{(x^2 - y^2)^2}{2} \leq \frac{s^2}{2} + z^2 s. \quad (\text{II.241.5})$$

Indeed $(x + y)^2 \leq 2s$. Summing, and using $\sum s^2 \leq (\sum s)^2$, gives the non-fixed part of (II.241.1). The fixed part is bounded by v_2^2 , proving both displayed inequalities. Now apply (II.239.7) and use $C_d + J_d = A_d$. Finally $E_4(r) \leq v^2$ and $T_4(r) \leq v^2$ imply (II.241.4). $\square$

Corollary II.241.2 (Profile-sensitive endpoint bound). *If $z \leq 1$, then*

$$\begin{aligned} \mathcal{U}_d(F) \leq & \left(1 - \frac{z}{2}\right)^q + A_d v_2^2 + \frac{C_d}{2}(v - v_2)^2 \\ & + C_d z^2(v - v_2) + B_d(v + a)v^2. \end{aligned} \quad (\text{II.241.6})$$

Proof. Use theorem II.240.3 and proposition II.241.1. The right side of (II.241.3) is nonnegative, and $0 < m^{q-4} \leq 1$. $\square$

For $v > 0$ write $\tau = v_2/v$. The leading quartic coefficient in (II.241.6) is

$$\Phi_d(\tau) = A_d \tau^2 + \frac{C_d}{2}(1 - \tau)^2. \quad (\text{II.241.7})$$

Since $A_d > C_d/2 > 0$ for $d \geq 2$, its maximum on $[0, 1]$ is A_d , attained only at $\tau = 1$. Thus the worse even-order coefficient is not merely a formal consequence of allowing torsion: it is attained only when the centered Fourier energy becomes concentrated in two-torsion.

II.242 The explicit local stability theorem

Source 55, delivered Section 7.

Define

$$K_{0,d} = \frac{q(q-1)}{8}, \quad (\text{II.242.1})$$

$$M_d = \begin{cases} C_d/2, & |G| \text{ odd}, \\ A_d, & |G| \text{ even}, \end{cases} \quad L_d = \begin{cases} C_d, & |G| \text{ odd}, \\ 0, & |G| \text{ even}, \end{cases} \quad (\text{II.242.2})$$

$$K_d = K_{0,d} + M_d, \quad H_d = B_d + L_d. \quad (\text{II.242.3})$$

For a radius and uniform inversion constant, put

$$K_* = K_{0,d} + A_d, \quad H_* = B_d + C_d, \quad c = q/2, \quad (\text{II.242.4})$$

$$\rho_d = \min \left\{ \frac{1}{2}, \frac{q}{16(K_* + 1)}, \sqrt{\frac{q}{8(B_d + 2C_d + 1)}} \right\}, \quad (\text{II.242.5})$$

$$R_d = \frac{12K_*^2}{c^5} + \frac{24K_*H_*}{c^6} + \frac{8H_*}{c^4}. \quad (\text{II.242.6})$$

All these constants are explicit. In (II.242.2), the even-order choice is valid on every group, though it is not optimal on odd-order groups.

Theorem II.242.1 (A posteriori local phase certificate). *Let $P \in \mathcal{P}_{d-1}(G)$ and rotate it so that $m = \mathbb{E}fP \geq 0$. Suppose $|f| \leq 1$ and*

$$z = 2(1 - m) = D + a \leq \rho_d, \quad D = \mathbb{E}|f - P|^2, \quad a = 1 - \mathbb{E}|f|^2. \quad (\text{II.242.7})$$

Then, writing $\varepsilon = 1 - \mathcal{U}_d(f)$,

$$\varepsilon \geq \frac{q}{2}D - K_d D^2 - H_d D^3 + \frac{q}{4}a, \quad (\text{II.242.8})$$

and

$$D + a \leq \frac{4}{q}\varepsilon. \quad (\text{II.242.9})$$

Moreover,

$$D \leq \frac{2}{q}\varepsilon + \frac{8K_d}{q^3}\varepsilon^2 + R_d\varepsilon^3. \quad (\text{II.242.10})$$

The phase P need not minimize distance.

Proof. Demodulate by P and use the notation of section II.238. Since $z \leq 1$, theorem II.240.3 and proposition II.241.1 give the master bound

$$\mathcal{U}_d(f) \leq \left(1 - \frac{z}{2}\right)^q + M_d D^2 + B_d z D^2 + L_d z^2 D. \quad (\text{II.242.11})$$

For odd order use $v_2 = 0$, $v \leq D$, and $v + a \leq z$. For even order use (II.241.4) instead.

Taylor's formula with a nonpositive third derivative gives

$$(1 - z/2)^q \leq 1 - \frac{q}{2}z + K_{0,d}z^2 \quad (0 \leq z \leq 1).$$

Substitute $z = D + a$ in (II.242.11) and collect the terms containing a to obtain

$$\begin{aligned} \varepsilon \geq & \frac{q}{2}D - K_d D^2 - H_d D^3 \\ & + a \left\{ \frac{q}{2} - K_{0,d}(2D + a) - B_d D^2 - L_d(2D + a)D \right\}. \end{aligned} \quad (\text{II.242.12})$$

Since $2D + a \leq 2\rho_d$, the expression in braces is at least

$$\frac{q}{2} - 2K_{0,d}\rho_d - (B_d + 2C_d)\rho_d^2 \geq \frac{q}{4}.$$

This proves (II.242.8). Also

$$K_d D + H_d D^2 \leq K_* \rho_d + H_* \rho_d^2 \leq \frac{q}{16} + \frac{q}{8} < \frac{q}{4},$$

which proves (II.242.9).

For clarity, we spell out the inversion rather than hiding it in an asymptotic notation. From (II.242.8), dropping the nonnegative amplitude term, and from (II.242.9),

$$D \leq \frac{\varepsilon}{c} + \frac{K_d}{c} D^2 + \frac{H_d}{c} D^3, \quad D \leq \frac{2\varepsilon}{c}.$$

Because $0 \leq \varepsilon \leq 1$, it follows first that

$$D - \frac{\varepsilon}{c} \leq C_1 \varepsilon^2, \quad C_1 = \frac{4K_d}{c^3} + \frac{8H_d}{c^4}.$$

Consequently (using the positive part if $D - \varepsilon/c < 0$)

$$D^2 \leq \frac{\varepsilon^2}{c^2} + \frac{3C_1}{c} \varepsilon^3, \quad D^3 \leq \frac{8\varepsilon^3}{c^3}.$$

Substituting once more yields

$$D \leq \frac{\varepsilon}{c} + \frac{K_d}{c^3} \varepsilon^2 + \left(\frac{12K_d^2}{c^5} + \frac{24K_d H_d}{c^6} + \frac{8H_d}{c^4} \right) \varepsilon^3.$$

Replacing K_d, H_d by K_*, H_* in the cubic coefficient proves (II.242.10). $\square$

The sharp coefficients simplify dramatically:

$$K_{0,d} + \frac{C_d}{2} = \frac{q(3^d - 1)}{16}, \quad K_{0,d} + A_d = \frac{q(q^2 - 1)}{24}. \quad (\text{II.242.13})$$

Thus $8K_d/q^3$ is exactly (II.237.2).

Corollary II.242.2 (Explicit U^3 certificate). *Suppose P is a quadratic phase in the difference sense, $|f| \leq 1$, and after constant rotation $2(1 - \mathbb{E}f\bar{P}) = D + a \leq 1/44$. With $\varepsilon = 1 - \|f\|_{U^3}^8$, one has*

$$\begin{aligned} \varepsilon & \geq 4D - 13D^2 - 620D^3 + 2a, & |G| \text{ odd}, \\ \varepsilon & \geq 4D - 21D^2 - 608D^3 + 2a, & |G| \text{ even}. \end{aligned}$$

In either case $D + a \leq \varepsilon/2$, and

$$D \leq \frac{\varepsilon}{4} + \begin{cases} 13\varepsilon^2/64, & |G| \text{ odd}, \\ 21\varepsilon^2/64, & |G| \text{ even} \end{cases} + \frac{25813}{256} \varepsilon^3.$$

Proof. Here $q = 8$, $C_3 = 12$, $A_3 = 14$, $B_3 = 608$, $K_* = 21$, $H_* = 620$, $\rho_3 = 1/44$, and (II.242.6) equals $25813/256$. $\square$

The radius $1/44$ is not asserted to be optimal. Unlike the quadratic coefficient, neither this radius nor the displayed cubic error constant has been optimized.

II.243 Uniform entry and matching extremizing families

Source 55, delivered Section 8.

II.243.1 The published entry theorem

Theorem II.243.1 (Eisner–Tao entry theorem, finite-group form). *For fixed d , if $|f| \leq 1$ and $\|f\|_{U^d} \rightarrow 1$, then there is $P \in \mathcal{P}_{d-1}(G)$ with $\mathbb{E}|f - P| = o_d(1)$, uniformly in the finite abelian group G .*

This is [7, Theorem 1.1], not a theorem newly proved here. Since $|f - P| \leq 2$, its conclusion implies squared L^2 distance $o_d(1)$. Also $a \leq \varepsilon$ by lemma II.238.1. For a minimizing phase, after rotation, $z = D_d(f) + a \rightarrow 0$ uniformly. Thus there is $e_d > 0$ such that $\varepsilon \leq e_d$ implies the hypothesis of theorem II.242.1. Equations (II.242.10) and (II.242.13) prove the upper bound in theorem II.237.1.

II.243.2 Separation of nonconstant phase classes

To prove sharpness, it is not enough to calculate distance to the constant phase: another polynomial phase could conceivably be closer. The following uniform gap excludes that possibility.

Lemma II.243.2 (A uniform phase-class gap). *If $R : G \rightarrow S^1$ is a phase polynomial of degree at most $s \geq 1$ and is not constant, then*

$$|\mathbb{E}R| \leq 1 - \gamma_s, \quad \gamma_s = 4^{-(s-1)}. \quad (\text{II.243.1})$$

Consequently, if $|F| = 1$ and

$$D_{\text{const}}(F) := \min_{c \in S^1} \mathbb{E}|F - c|^2 < \gamma_{d-1}/2,$$

then $D_d(F) = D_{\text{const}}(F)$.

Proof. For $s = 1$, a phase is a constant times a character, and a nonconstant character has mean zero. Suppose the assertion known for $s - 1$. If $|\mathbb{E}R| > 1 - \gamma_s$, rotate R so that its mean is positive, and write $\eta = 1 - \mathbb{E}R < \gamma_s$. Then $\|R - 1\|_2^2 = 2\eta$. For each direction h , the phase derivative $\Delta_h R(x) = R(x + h)\overline{R(x)}$ has degree at most $s - 1$, and

$$\|\Delta_h R - 1\|_2 \leq 2\|R - 1\|_2.$$

It follows that $\text{Re } \mathbb{E}\Delta_h R \geq 1 - 4\eta > 1 - \gamma_{s-1}$. By induction every derivative is constant. Therefore R is a constant multiple of a character. Its nonzero mean forces that character to be trivial, a contradiction. This proves (II.243.1).

For the final assertion choose a closest constant c . If a nonconstant phase P had distance to F at most $\sqrt{D_{\text{const}}(F)}$, the triangle inequality would give

$$\|P - c\|_2 \leq 2\sqrt{D_{\text{const}}(F)} < \sqrt{2\gamma_{d-1}}.$$

But (II.243.1) implies $\min_c \|P - c\|_2^2 = 2(1 - |\mathbb{E}P|) \geq 2\gamma_{d-1}$. $\square$

II.243.3 Even-order groups: an exact one-character family

Every even-order finite abelian group has a nontrivial character $\chi : G \rightarrow \{-1, 1\}$. For t real put

$$F_t = e^{it\chi}. \quad (\text{II.243.2})$$

Its mean is $\cos t$. For sufficiently small t , lemma II.243.2 gives

$$D_d(F_t) = 2(1 - \cos t). \quad (\text{II.243.3})$$

The additive cubical difference of χ factors exactly:

$$\sum_{\omega} (-1)^{|\omega|} \chi(x + \omega \cdot h) = \chi(x) \prod_{j=1}^d (1 - \chi(h_j)).$$

It vanishes unless all $\chi(h_j) = -1$, an event of probability $1/q$; on that event it is $q\chi(x)$. Hence

$$\varepsilon(F_t) = \frac{1 - \cos(qt)}{q}. \quad (\text{II.243.4})$$

For small positive e , choose $t = q^{-1} \arccos(1 - qe)$. This gives the exact attained distance

$$\mathcal{D}_d^{\text{even}}(e) = 2 \left[1 - \cos \left(\frac{1}{q} \arccos(1 - qe) \right) \right]. \quad (\text{II.243.5})$$

Taylor expansion, or elimination of t^2 from (II.243.3)–(II.243.4), gives

$$\mathcal{D}_d^{\text{even}}(e) = \frac{2}{q}e + \frac{q^2 - 1}{3q^2}e^2 + O_d(e^3).$$

The choice of t and the phase-gap range are independent of the group.

II.243.4 Odd-order groups: a conjugate-pair family

Let G have odd order and be nontrivial. Choose a nontrivial character χ , so $\chi \neq -\chi$ in additive dual notation, and set

$$b = \sqrt{2} \operatorname{Re}(\zeta\chi), \quad F_t = e^{itb}, \quad |\zeta| = 1. \quad (\text{II.243.6})$$

Then $\mathbb{E}b = 0$, $\mathbb{E}b^2 = 1$, $|b| \leq \sqrt{2}$, and the Fourier mass of b is $1/2$ at each of $\chi, -\chi$. Thus

$$Q_d(b) = C_d/2. \quad (\text{II.243.7})$$

Write $\mu_4 = \mathbb{E}b^4$ and $S = \sum_{\omega} (-1)^{|\omega|} b(X_{\omega})$. Triple independence gives

$$\mathbb{E}S^2 = q, \quad (\text{II.243.8})$$

$$\mathbb{E}S^4 = q\mu_4 + 3q(q - 1) + 24Q_d(b). \quad (\text{II.243.9})$$

Indeed the multiplicity patterns $3 + 1$ and $2 + 1 + 1$ vanish, the pattern $2 + 2$ contributes $6\binom{q}{2}$, and the four-distinct pattern contributes the final term. Its signs are positive on every dependent four-set by proposition II.239.3.

The law of S is symmetric: interchanging the two endpoints of the first cube direction changes S to $-S$ and preserves the sampling measure. Consequently

$$\varepsilon(F_t) = \frac{q}{2}t^2 - \left(\frac{q\mu_4}{24} + \frac{q(q-1)}{8} + \frac{C_d}{2} \right) t^4 + O_d(t^6). \quad (\text{II.243.10})$$

For sufficiently small t , the closest phase is a constant by lemma II.243.2, and

$$D_d(F_t) = 2(1 - |\mathbb{E}e^{itb}|) = t^2 - \frac{\mu_4}{12}t^4 + O(t^6). \quad (\text{II.243.11})$$

To check the latter formula, expand the mean through order four. Its real part is $1 - t^2/2 + \mu_4 t^4/24 + O(t^6)$, while its imaginary part is $O(t^3)$ because $\mathbb{E}b = 0$. Taking the modulus changes it only at order six. All error estimates are uniform because $|b| \leq \sqrt{2}$.

Eliminating t^2 makes the fourth moment cancel:

$$\varepsilon(F_t) = \frac{q}{2}D_d(F_t) - \left(K_{0,d} + \frac{C_d}{2} \right) D_d(F_t)^2 + O_d(D_d(F_t)^3).$$

Inversion gives the odd coefficient in (II.237.2). The derivative of (II.243.10) is $qt + O_d(t^3)$, uniformly, so every sufficiently small positive defect is attained by a positive t . Together with the uniform upper bound, the two families prove (II.237.3) and complete the proof of theorem II.237.1.

II.244 Proof of second-order rigidity

Source 55, delivered Section 9.

We now prove theorem II.237.3. All little-oh estimates in this section are along its given sequence. Demodulate by a minimizing phase and use F, r, D, a, z, v, v_2 as before, suppressing the index n . Write $c = q/2$, and let $M = M_d$, $K = K_{0,d} + M$.

II.244.1 Amplitude and quartic saturation

By theorem II.237.1 and the saturation hypothesis,

$$D = \frac{\varepsilon}{c} + \frac{K}{c^3}\varepsilon^2 + o(\varepsilon^2), \quad cD - KD^2 = \varepsilon + o(\varepsilon^2). \quad (\text{II.244.1})$$

The local theorem gives $z = O_d(\varepsilon)$ and

$$\varepsilon \geq cD - KD^2 - H_d D^3 + \frac{q}{4}a.$$

Since $a \geq 0$, (II.244.1) implies

$$a = o(\varepsilon^2), \quad v = D - z^2/4 \sim D \sim \varepsilon/c. \quad (\text{II.244.2})$$

The exact expansion (II.240.6), expanded in z , now yields

$$\varepsilon = cD - K_{0,d}D^2 - \text{Re } Q_d(r) + o(\varepsilon^2).$$

Combining with (II.244.1),

$$\text{Re } Q_d(r) = Mv^2 + o(v^2). \quad (\text{II.244.3})$$

II.244.2 Even order: concentration at one real character

In even order, (II.241.3) and (II.244.3) imply

$$A_d \leq \Phi_d(v_2/v) + o(1).$$

The unique maximum of (II.241.7) is at 1, so $v_2/v \rightarrow 1$. The non-fixed Fourier mass is $o(v)$; in particular its fourth-power sum is $o(v^2)$. Thus

$$E_4(r) = T_4(r) + o(v^2).$$

Applying (II.239.7) once more to (II.244.3) gives $T_4(r) \sim v^2$. Let the squared Fourier masses on nontrivial order-two characters be p_χ . Since

$$\sum p_\chi \sim v, \quad \sum p_\chi^2 \sim v^2, \quad \sum p_\chi^2 \leq (\max p_\chi) \sum p_\chi,$$

one such character χ carries $(1 - o(1))v$. Equation (II.238.4) at a fixed character gives $|2 \operatorname{Re} \hat{r}(\chi)| \leq z = O(\varepsilon) = o(\sqrt{v})$. Therefore for a sign s ,

$$\hat{r}(\chi) = is\sqrt{v} + o(\sqrt{v}), \quad \|r - is\sqrt{v}\chi\|_2 = o(\sqrt{v}).$$

Finally $m - 1 = O(D) = o(\sqrt{D})$ and $v/D \rightarrow 1$, which proves (II.237.5).

II.244.3 Odd order: concentration at one conjugate pair

There is no nontrivial order-two character. Since $C_d > 0$, (II.239.7), (II.241.1), and (II.244.3) give

$$E_4(r) = \frac{1}{2}v^2 + o(v^2).$$

Let s_j be the total squared mass in the j th conjugate pair. The pairwise inequality (II.241.5) implies

$$E_4(r) \leq \frac{1}{2} \sum_j s_j^2 + z^2 v.$$

Here $z^2 v = o(v^2)$, while $\sum_j s_j = v$. Thus $\sum_j s_j^2 \sim v^2$, and one pair $\{\chi, -\chi\}$ carries $(1 - o(1))v$. For its coefficients $\alpha = \hat{r}(\chi)$ and $\beta = \hat{r}(-\chi)$, (II.238.4) says

$$\beta = -\bar{\alpha} + o(\sqrt{v}), \quad |\alpha| = |\beta| + o(\sqrt{v}) = \sqrt{v/2} + o(\sqrt{v}).$$

Choose $\zeta \in S^1$ so that $\alpha = i\sqrt{v/2}\zeta + o(\sqrt{v})$. Fourier inversion and Parseval give

$$\|r - i\sqrt{2v} \operatorname{Re}(\zeta\chi)\|_2 = o(\sqrt{v}).$$

As above, adding back $m - 1$ and replacing v by D proves (II.237.6). This completes the proof of theorem II.237.3.

II.245 An exact calibration at degree two

Source 55, delivered Section 10.

At $d = 2$ the Fourier identity allows a nonasymptotic answer on every even-order group. This elementary calibration is useful for checking the normalizations and the direction of the inequality; it is not being claimed as an independent historical novelty.

Proposition II.245.1 (Exact even-order U^2 stability modulus). *If G is a nontrivial finite abelian group of even order, then for $0 \leq e \leq 1/2$,*

$$\mathfrak{S}_{2,G}(e) = 2 \left(1 - \sqrt{\frac{1 + \sqrt{1 - 2e}}{2}} \right). \quad (\text{II.245.1})$$

The right side remains a valid upper bound for every finite abelian group.

Proof. Polynomial phases of degree one are constants times characters. Let $s = \mathbb{E}|f|^2$ and $m = \max_{\chi} |\hat{f}(\chi)|$. Then

$$D_2(f) = 1 + s - 2m, \quad \mathcal{U}_d(f) = \sum_{\chi} |\hat{f}(\chi)|^4 \leq m^2 s \leq m^2.$$

If $\varepsilon(f) \leq e \leq 1/2$, it follows that $m \geq \sqrt{1 - e}$ and $D_2(f) \leq 2 - \sqrt{2}$. Demodulate by a maximizing character. With $D = D_2(f)$ and $z = D + a$,

$$\mathcal{U}_d(f) \leq m^4 + v^2 = \left(1 - \frac{z}{2}\right)^4 + \left(D - \frac{z^2}{4}\right)^2.$$

For fixed D , this expression is decreasing in z from D to the actual value of z : its derivative is

$$-2(1 - z/2)^3 - z(D - z^2/4) \leq 0.$$

Both terms have the indicated sign on that interval because $m \geq 0$ and the endpoint variance is nonnegative. Consequently

$$\varepsilon(f) \geq g(D) := 2D - \frac{5}{2}D^2 + D^3 - \frac{1}{8}D^4.$$

Writing $u = 1 - D/2$, this is $g(D) = 2u^2(1 - u^2)$. The function g increases on $0 \leq D \leq 2 - \sqrt{2}$, so $g(D) \leq e$ implies the upper bound in (II.245.1).

For equality choose an order-two character and $f = e^{it\chi} = \cos t + i\chi \sin t$, with $0 \leq t \leq \pi/4$. The largest Fourier coefficient is $\cos t$, so the closest phase is the constant one. Its defect is $(1 - \cos 4t)/4$, covering $[0, 1/2]$, and its distance is $2(1 - \cos t)$. This is exactly (II.245.1). $\square$

In particular, the even-order example curve (II.243.5) is exactly optimal at $d = 2$. Our higher-degree theorem proves optimality through second order, not the full exact curve for $d \geq 3$.

II.246 Consequences for phase testing and local structure

Source 55, delivered Section 11.

II.246.1 A sharp soundness expansion for unit-modulus cube tests

Suppose $|f| = 1$ and define the squared cubical test loss

$$\mathcal{L}_d(f) = \mathbb{E} \left| \prod_{\omega} \mathcal{C}^{|\omega|} f(X_{\omega}) - 1 \right|^2.$$

Since the product has modulus one and its expectation is real, $\mathcal{L}_d(f) = 2\varepsilon(f)$. Therefore theorem II.237.1 gives

$$D_d(f) \leq \frac{1}{q} \mathcal{L}_d(f) + \frac{b_d(G)}{4} \mathcal{L}_d(f)^2 + O_d(\mathcal{L}_d(f)^3). \quad (\text{II.246.1})$$

Both coefficients are sharp by the same unit-modulus examples. This is an average squared-loss statement for complex phases, not a claim about the rejection probability of a Boolean-valued test without a conversion of the loss convention.

II.246.2 Certification after a coarse inverse step

Suppose a separate argument supplies a polynomial phase P with $|\mathbb{E} f \overline{P}| \geq 1 - \rho_d/2$. Rotate P and compute the cubical loss. The local theorem immediately returns an explicit bound on $\mathbb{E}|f - P|^2$, without needing the qualitative entry theorem or a new minimization over all phases. At $d = 3$, the candidate-correlation condition is simply $|\mathbb{E} f \overline{P}| \geq 1 - 1/88$.

This supplies a concrete endpoint interface for the extraction and integration architecture in Gowers's proof. It does not supply such a candidate from an arbitrary fixed positive U^d norm. In particular it should not be substituted for an open Section 16 or Section 17 catalogue entry whose hypotheses are substantially weaker.

II.246.3 Cosets and quotients

If the same problem is posed on a coset $x_0 + H$ with normalized counting measure and all directions in H , translation identifies it exactly with the theorem on H . The parity in the coefficient is the parity of $|H|$, not that of an unrelated ambient group. Pullback along a surjective homomorphism preserves every cube average and the explicit example families. For the upper theorem no quotient reduction is necessary.

The corresponding assertion for an arbitrary Bohr set with restricted sampling is not automatic: the proof depends on exact triple independence. Likewise, restricting all vertices to an interval changes the underlying sampling law. These are substantive extension questions, not harmless changes of notation.

II.247 Near-maximal uniformity and global frequency maps

Source 57, delivered Section 5.

The small interval near $\alpha = 1$ in the preceding argument admits a different, much stronger conclusion: the progressions can be the whole prime field. We prove this directly, with explicit constants, by correcting the dominant Fourier frequencies of multiplicative derivatives. The same argument extends to every uniformity order. On a prime field, permutation symmetry reduces the dependence on the number of derivative parameters to a linear loss.

All averages and Fourier transforms in this section are normalized. Thus, on $\mathbb{F}_p$, $\hat{u}(r) = \mathbb{E}_x u(x) e_p(-rx)$, where $e_p(t) = \exp(2\pi it/p)$; the Fourier coefficients in Gowers's original normalization are $p\hat{u}(r)$. We write

$$Q_j(u) = \|u\|_{U^j}^{2^j}, \quad \Delta_h u(x) = u(x+h) \overline{u(x)}.$$

On a general finite abelian group, frequencies belong to its character group, written additively. Distances between finite-valued maps mean the proportion of inputs on which they disagree. The near-extremizer principle that very large uniformity forces closeness to a polynomial phase is established in much greater generality by Eisner and Tao [7, Theorem 1.1]. The contribution here is an explicit frequency-correction argument suited to the preceding quantitative application, its symmetry improvement, and the constants in the resulting prime-field estimates.

II.247.1 Phase normalization and additive correction

Lemma II.247.1 (Normalization without an order-dependent loss). *Let G be finite abelian, $d \geq 1$, and $|f| \leq 1$. Define $g = f/|f|$ where $f \neq 0$, and $g = 1$ at the zeros of f . If $Q_d(f) \geq 1 - \varepsilon$, then*

$$\|f - g\|_1 \leq \varepsilon, \quad Q_d(g) \geq 1 - 2\varepsilon. \quad (\text{II.247.1})$$

Furthermore, for every $s \geq 1$ and every choice of directions,

$$\|\Delta_{h_1} \cdots \Delta_{h_s} f - \Delta_{h_1} \cdots \Delta_{h_s} g\|_1 \leq 2^s \varepsilon. \quad (\text{II.247.2})$$

Proof. For each cube, write its alternating product for f as AC , where A is the product of the 2^d absolute values and C is the corresponding product for g . Thus $0 \leq A \leq 1$ and $|C| = 1$. For $-1 \leq r \leq 1$,

$$2(1 - Ar) - (1 - r) = 1 + (1 - 2A)r \geq 1 - |1 - 2A| \geq 0.$$

Apply this with $r = \Re C$ and average. Since the cube averages are real, $1 - Q_d(g) \leq 2(1 - Q_d(f))$. Also $Q_d(f) \leq \mathbb{E}A \leq \mathbb{E}|f|$, using any fixed cube vertex, and consequently $\|f - g\|_1 = \mathbb{E}(1 - |f|) \leq \varepsilon$. Finally telescope the 2^s factors in a multiplicative derivative. Every other factor has modulus at most one, and translation preserves the L^1 norm. This proves (II.247.2) uniformly in the directions. $\square$

The next argument is a direct form of the classical Blum–Luby–Rubinfeld correction method [49]; we include the finite abelian group proof and constants.

Lemma II.247.2 (An elementary additive correction bound). *Let A, H be finite abelian groups and $u : A \rightarrow H$. Put*

$$\theta = \mathbb{P}_{x,y}[u(x+y) \neq u(x) + u(y)].$$

If $\theta < 1/6$, there is a homomorphism $v : A \rightarrow H$ such that

$$\text{dist}(u, v) \leq \frac{\theta}{1 - 2\theta} \leq 2\theta. \quad (\text{II.247.3})$$

Proof. Fix x and put $A_x(y) = u(x+y) - u(y)$. For independent uniform y, z , the addition tests at $(x+y, z-y)$ and $(y, z-y)$ each fail with probability θ : each ordered pair is uniformly distributed on A^2 . If both tests hold, then $A_x(y) = A_x(z)$. Thus the collision probability of A_x is at least $1 - 2\theta$. For any finite probability distribution (m_a) , $\sum_a m_a^2 \leq \max_a m_a$. The most frequent value of A_x therefore has mass at least $1 - 2\theta > 1/2$. Denote this unique value by $v(x)$.

The exact identity

$$A_{x+x'}(y) = A_x(x' + y) + A_{x'}(y)$$

and the union bound show that all three majority values occur simultaneously with probability at least $1 - 6\theta > 0$. Hence $v(x + x') = v(x) + v(x')$. If $u(x) \neq v(x)$, at least $1 - 2\theta$ of the original tests with first argument x fail. Writing $\delta = \text{dist}(u, v)$ gives $\delta(1 - 2\theta) \leq \theta$, as required. $\square$

Lemma II.247.3 (Symmetric bilinear correction on a prime field). *Let $\phi : \mathbb{F}_p^2 \rightarrow \mathbb{F}_p$ be symmetric, with $\phi(h, 0) = 0$. If*

$$\mathbb{P}_{h,h',k}[\phi(h + h', k) \neq \phi(h, k) + \phi(h', k)] \leq \tau, \quad (\text{II.247.4})$$

then some $c \in \mathbb{F}_p$ satisfies

$$\mathbb{P}_{h,k}[\phi(h, k) \neq chk] \leq 18\tau. \quad (\text{II.247.5})$$

No smallness condition on τ is required.

Proof. For each k , let θ_k be the failure rate in the row $h \mapsto \phi(h, k)$. If $\theta_k < 1/6$, apply Lemma II.247.2; the corrected row is $h \mapsto a(k)h$ and has distance at most $2\theta_k$. Otherwise take $a(k) = 0$, incurring distance at most $1 \leq 6\theta_k$. The zero row is already correct, so take $a(0) = 0$. The global distance δ from ϕ to $L(h, k) = a(k)h$ is at most 6τ .

Symmetry gives $D := \mathbb{P}[L(h, k) \neq L(k, h)] \leq 2\delta$. Put $q = (p - 1)/p$ and

$$m_c = \mathbb{P}[a(k)/k = c \mid k \neq 0].$$

Because $a(0) = 0$, a mismatch requires $h, k \neq 0$, and therefore

$$D = q^2 \left(1 - \sum_c m_c^2 \right).$$

Choose c with maximum m_c . The distance between L and $(h, k) \mapsto chk$ is $q^2(1 - \max_c m_c) \leq D$. Consequently the distance from ϕ to this bilinear map is at most $\delta + D \leq 3\delta \leq 18\tau$. The calculation is valid also for $p = 2$. $\square$

II.247.2 From Fourier concentration to addition tests

Lemma II.247.4 (Dominant frequencies of mixed derivatives). *Let G be finite abelian, $s \geq 1$, and $g : G \rightarrow \mathbb{C}$ be unimodular. Suppose $Q_{s+2}(g) \geq 1 - \rho$. There is a frequency map $\phi : G^s \rightarrow \widehat{G}$ with the following properties. It is symmetric in its arguments and equals zero whenever one argument is zero. There is a function $t : G^s \rightarrow [0, 1]$ with $\mathbb{E}t \leq \rho$ such that*

$$\left| \widehat{\Delta_{h_1} \cdots \Delta_{h_s} g}(\phi(h_1, \dots, h_s)) \right|^2 \geq 1 - t(h_1, \dots, h_s). \quad (\text{II.247.6})$$

In each coordinate, the additive-test failure rate of ϕ is at most 9ρ .

Proof. Write $F_{\mathbf{h}} = \Delta_{h_1} \cdots \Delta_{h_s} g$ and $t(\mathbf{h}) = 1 - Q_2(F_{\mathbf{h}})$. The Gowers recursion gives $\mathbb{E}t \leq \rho$. Since $F_{\mathbf{h}}$ is unimodular, Parseval and the Fourier expression for Q_2 give

$$\sum_r |\widehat{F_{\mathbf{h}}}(r)|^2 = 1, \quad \sum_r |\widehat{F_{\mathbf{h}}}(r)|^4 = 1 - t(\mathbf{h}).$$

Hence the largest squared Fourier modulus is at least $1 - t$. Choose a maximizing frequency with one fixed ordering of $\widehat{G}$ to break ties. Mixed derivatives commute, so this choice is symmetric. If one direction is zero, the derivative is identically one, and the maximizing frequency is uniquely zero.

For suitable unit scalars z_h , put

$$e(\mathbf{h}) = \left\| F_h - z_h \chi_{\phi(\mathbf{h})} \right\|_2^2 = 2 - 2|\widehat{F_h}(\phi(\mathbf{h}))| \leq 2t(\mathbf{h}).$$

Thus $\mathbb{E}e \leq 2\rho$. In the first coordinate the exact cocycle identity is

$$F_{h+h', \mathbf{k}}(x) = F_{h, \mathbf{k}}(x + h') F_{h', \mathbf{k}}(x).$$

If the selected frequencies fail to add, the unit characters approximating the two sides are orthogonal, regardless of their scalar phases. Their L^2 distance is $\sqrt{2}$. The triangle inequality, followed by Cauchy–Schwarz, therefore yields

$$\mathbf{1}_{\text{failure}} \leq \frac{3}{2}(e(h + h', \mathbf{k}) + e(h, \mathbf{k}) + e(h', \mathbf{k})).$$

Averaging bounds the failure rate by $(9/2)\mathbb{E}e \leq 9\rho$. The other coordinates are identical. $\square$

Theorem II.247.5 (A global square near maximal fourth uniformity). *Let p be prime and $f : \mathbb{F}_p \rightarrow \mathbb{C}$ satisfy $|f| \leq 1$ and*

$$Q_4(f) \geq 1 - \varepsilon, \quad 0 \leq \varepsilon \leq 1/1024.$$

There are $c \in \mathbb{F}_p$ and $B \subseteq \mathbb{F}_p^2$ such that

$$|B| \geq (1 - 332\varepsilon)p^2, \quad \left| \widehat{\Delta_h \Delta_k f}(chk) \right| \geq \frac{3}{4} - 4\varepsilon \quad ((h, k) \in B). \quad (\text{II.247.7})$$

In particular $|B| \geq (173/256)p^2 > p^2/2$, and every displayed Fourier coefficient is at least $191/256 > 1/2$.

Proof. Normalize f to g by Lemma II.247.1. Apply Lemma II.247.4 with $s = 2$ and $\rho = 2\varepsilon$. The first-coordinate test fails with probability at most $\tau = 18\varepsilon$, so Lemma II.247.3 supplies c with $\phi(h, k) = chk$ except on at most 324ε of the pairs. Discard also the pairs where $t(h, k) > 1/4$; Markov’s inequality bounds their proportion by $4\mathbb{E}t \leq 8\varepsilon$. On each remaining pair, (II.247.6) gives a Fourier modulus at least $\sqrt{3}/2 \geq 3/4$ for $\Delta_h \Delta_k g$. The uniform bound (II.247.2) changes this coefficient by at most 4ε on returning to f . The exceptional proportions sum to 332ε . At $\varepsilon = 1/1024$, the two elementary estimates in the statement are equal to $173/256$ and $191/256$. $\square$

[write] **7 October 2026; sources 54–59.** *Formalized since the pin*, as the repository’s own proof of this theorem (`endpoint_global_fourier`, `Proofs13EndpointGlobalFourier.lean`, line 12, `d48e8af00`: density at least $1 - 332\varepsilon$ and coefficient at least $(3/4 - 4\varepsilon)N$ in the unnormalized convention, with no threshold and $\varepsilon = 0$ included). Its bilinear correction (`endpoint_symmetric_rank_one`, `Proofs13EndpointSymmetricRankOne.lean`, line 11) is a pivot argument that does not need the zero row $\phi(h, 0) = 0$ assumed in Lemma II.247.3, as `FORMALIZATION_STATUS.txt` records. Checked at the write: $1 - 332/1024 = 173/256$ and $3/4 - 4/1024 = 191/256$.

This theorem supplies the endpoint needed in Section 13 with $P = Q = \mathbb{F}_p$, represented by $0, 1, \dots, p - 1$ with common step one. It uses neither a scale threshold nor an inverse theorem as an input. The exact endpoint $\varepsilon = 0$ is included.

II.247.3 Every order on a prime field

Lemma II.247.6 (Correction of a symmetric table in several variables). *Let $\phi : \mathbb{F}_p^s \rightarrow \mathbb{F}_p$ be symmetric and vanish whenever any coordinate is zero. If its first-coordinate additive-test failure rate is at most τ , then some $c \in \mathbb{F}_p$ satisfies*

$$\mathbb{P}[\phi(h_1, \dots, h_s) \neq ch_1 \cdots h_s] \leq 12s\tau. \quad (\text{II.247.8})$$

There is no smallness condition on τ .

Proof. Correct the first-coordinate rows as in Lemma II.247.3. This gives $L(\mathbf{h}) = a(h_2, \dots, h_s)h_1$ at global distance $\delta \leq 6\tau$. Choose $a = 0$ on every row for which one remaining coordinate is zero; those original rows are identically zero. Both L and ϕ vanish when $h_1 = 0$ as well. Thus every discrepancy lies in $\Omega = (\mathbb{F}_p^*)^s$, of density q^s , where $q = (p-1)/p$.

On Ω define

$$\psi(\mathbf{h}) = \frac{\phi(\mathbf{h})}{h_1 \cdots h_s}.$$

It is symmetric and is within δ/q^s of the predictor $a(h_2, \dots, h_s)/(h_2 \cdots h_s)$, which is independent of h_1 . Symmetry gives, for each coordinate j , a predictor independent of h_j with the same error bound.

Take independent uniform $X, Y \in \Omega$ and replace the coordinates of X by those of Y successively. At replacement j , the corresponding predictor agrees on the two tuples. Their ψ values can differ only if at least one prediction is wrong, with probability at most $2\delta/q^s$. Hence

$$\mathbb{P}[\psi(X) \neq \psi(Y)] \leq \frac{2s\delta}{q^s}.$$

For the distribution $m_c = \mathbb{P}[\psi = c]$, the left side is $1 - \sum_c m_c^2$. A most frequent c has exceptional probability $1 - \max_c m_c \leq 1 - \sum_c m_c^2$. Multiplying the resulting estimate by q^s and including the exactly correct zero-coordinate tuples gives (II.247.8). $\square$

Theorem II.247.7 (Symmetric frequency extraction at every order). *Let p be prime, $s \geq 1$, and $f : \mathbb{F}_p \rightarrow \mathbb{C}$ satisfy $|f| \leq 1$ and $Q_{s+2}(f) \geq 1 - \varepsilon$. There exist $c \in \mathbb{F}_p$ and $B \subseteq \mathbb{F}_p^s$ such that*

$$|B| \geq (1 - (216s + 8)\varepsilon)p^s, \quad (\text{II.247.9})$$

$$\left| \widehat{\Delta_{h_1} \cdots \Delta_{h_s} f}(ch_1 \cdots h_s) \right| \geq \frac{3}{4} - 2^s \varepsilon \quad (\mathbf{h} \in B). \quad (\text{II.247.10})$$

A half-density set with Fourier modulus at least $1/2$ follows whenever

$$\varepsilon \leq \min \left\{ \frac{1}{432s + 16}, 2^{-s-2} \right\}.$$

For $s = 2$, Theorem II.247.5 improves the coefficient 440 in (II.247.9) to 332.

Proof. After phase normalization, Lemma II.247.4 has $\rho = 2\varepsilon$ and test error at most 18ε . Lemma II.247.6 corrects the selector outside a set of proportion $216s\varepsilon$. Removing $t > 1/4$ costs a further 8ε . On the retained set the Fourier modulus for the phase is at least $3/4$; (II.247.2) gives the stated loss $2^s \varepsilon$ for f . $\square$

II.247.4 An extension to arbitrary finite abelian groups

The scalar division used above is specific to a one-dimensional prime field. A separate additive-correction argument works on all finite abelian groups, with a larger dependence on s .

Lemma II.247.8 (Correction of an approximately separately additive map). *Let $A_1, \dots, A_s, H$ be finite abelian groups. Suppose $\phi : A_1 \times \dots \times A_s \rightarrow H$ has additive-test failure at most τ in every coordinate. Put*

$$C_1 = 2, \quad C_s = 6 + 38C_{s-1} = \frac{80 \cdot 38^{s-1} - 6}{37}. \quad (\text{II.247.11})$$

If $38^{s-1}\tau < 1/6$, there is a separately additive map B such that $\text{dist}(\phi, B) \leq C_s\tau$.

Proof. The case $s = 1$ is Lemma II.247.2. For larger s , correct the first-coordinate rows with global distance at most 6τ , using the zero homomorphism on rows whose test error is at least $1/6$. The corrected rows form a map a of the remaining coordinates into the finite abelian group $\text{Hom}(A_1, H)$.

In any remaining coordinate, the corrected addition test, after evaluation at a uniform element of A_1 , can fail only if the original test fails or one of its three values was changed. Its failure probability is therefore at most $\tau + 18\tau = 19\tau$. A nonzero homomorphism $A_1 \rightarrow H$ has a proper kernel, so it is nonzero on at least half of A_1 . The actual homomorphism-valued addition test for a consequently fails with probability at most 38τ . Apply induction in $s - 1$ coordinates with this parameter. Its premise is precisely $38^{s-1}\tau < 1/6$, and the total correction distance is at most $6\tau + 38C_{s-1}\tau = C_s\tau$. Evaluating the resulting homomorphism-valued map at the first coordinate produces the required separately additive B . $\square$

Theorem II.247.9 (Frequency extraction on a finite abelian group). *Let G be finite abelian, $s \geq 1$, and $f : G \rightarrow \mathbb{C}$ satisfy $|f| \leq 1$ and $Q_{s+2}(f) \geq 1 - \varepsilon$. If $18 \cdot 38^{s-1}\varepsilon < 1/6$, there is an s -additive map $B : G^s \rightarrow \widehat{G}$ and a set $E \subseteq G^s$ of density at least $1 - (18C_s + 8)\varepsilon$ such that*

$$\left| \widehat{\Delta_{h_1} \cdots \Delta_{h_s} f(B(\mathbf{h}))} \right| \geq \frac{3}{4} - 2^s \varepsilon \quad (\mathbf{h} \in E).$$

Proof. Combine Lemmas II.247.1, II.247.4, and II.247.8 with $\tau = 18\varepsilon$. The selector-correction error is $18C_s\varepsilon$; removing $t > 1/4$ costs 8ε , and transfer back to f costs $2^s\varepsilon$ in the Fourier modulus. $\square$

Theorem II.247.9 asserts separate additivity of a frequency map. It does not assert that this map is a symmetric tensor or that it integrates to one polynomial phase. Such integration requires an additional argument. The next result carries it out in the cubic prime-field case.

II.247.5 An explicit cubic inverse theorem near the maximum

Corollary II.247.10 (Cubic phase approximation with an explicit error). *Let $p \geq 5$ be prime and $f : \mathbb{F}_p \rightarrow \mathbb{C}$ satisfy $|f| \leq 1$ and*

$$Q_4(f) \geq 1 - \varepsilon, \quad 0 \leq \varepsilon < 1/17604.$$

There is $P \in \mathbb{F}_p[X]$, of degree at most three, such that

$$|\mathbb{E}_x f(x) e_p(-P(x))| \geq \sqrt{1 - 6194\varepsilon} - \varepsilon. \quad (\text{II.247.12})$$

For a suitable scalar z of modulus one,

$$\|f - z e_p(P)\|_2^2 \leq 2(1 - \sqrt{1 - 6194\varepsilon}) + 2\varepsilon \leq 12390\varepsilon. \quad (\text{II.247.13})$$

Proof. We first record the degree descent for a unimodular function u with $Q_4(u) \geq 1 - D$. Lemma II.247.4 gives mean deficit at most D and first-coordinate test failure at most $9D$. The symmetric bilinear correction changes at most $162D$ of its frequencies. Therefore, using (II.247.6) on the agreement set and positivity on its complement, some c satisfies

$$\mathbb{E}_{h,k} \left| \widehat{\Delta_h \Delta_k u}(chk) \right|^2 \geq 1 - 163D. \quad (\text{II.247.14})$$

There is no thresholding of the individual deficits here.

Set $u_1(x) = u(x)e_p(-(c/6)x^3)$. The additive polynomial identity

$$((x+h+k)^3 - (x+h)^3 - (x+k)^3 + x^3)c/6 = chkx + (c/2)hk(h+k)$$

shows that the squared modulus of $\mathbb{E}_x \Delta_h \Delta_k u_1(x)$ equals the squared Fourier modulus in (II.247.14). The Gowers recursion consequently gives $Q_3(u_1) \geq 1 - 163D$.

For a unimodular v with $Q_3(v) \geq 1 - D_1$, the one-variable selector of Lemma II.247.4 has addition-test error at most $9D_1$. If $9D_1 < 1/6$, Lemma II.247.2 supplies a frequency bh outside a set of proportion at most $18D_1$. The same averaging gives

$$\mathbb{E}_h \left| \widehat{\Delta_h v}(bh) \right|^2 \geq 1 - 19D_1.$$

Demodulate by $e_p(-(b/2)x^2)$. Since $((x+h)^2 - x^2)b/2 = bhx + (b/2)h^2$, the resulting function v_1 satisfies $Q_2(v_1) \geq 1 - 19D_1$. Parseval gives a frequency r for which $|\widehat{v}_1(r)|^2 \geq 1 - 19D_1$.

Now normalize the original f to its unit phase g . Lemma II.247.1 permits $D = 2\varepsilon$, whence $D_1 = 326\varepsilon$. The last correction requires $2934\varepsilon < 1/6$, exactly the stated bound $\varepsilon < 1/17604$, and $19D_1 = 6194\varepsilon$. Undoing the demodulations gives

$$P(X) = (c/6)X^3 + (b/2)X^2 + rX, \quad |\mathbb{E}_g e_p(-P)| \geq \sqrt{1 - 6194\varepsilon}.$$

Both divisions are legitimate because $p \geq 5$. Since $\|f - g\|_1 \leq \varepsilon$, this proves (II.247.12). Align a unit scalar z with the correlation and use $\mathbb{E}|f|^2 \leq 1$ to obtain the first bound in (II.247.13). The second uses $1 - \sqrt{1-x} \leq x$ for $0 \leq x \leq 1$. $\square$

The scalar z is necessary: a constant unit phase need not be a p th root of unity. This corollary is an explicit version of a known near-extremizer phenomenon, not a new qualitative inverse principle. Eisner and Tao explicitly leave their effective polynomial rates unspecified [7, Remark 1.6]. No optimality is claimed for the constants 6194 and 12390.

II.247.6 Explicit integration in every degree at higher characteristic

The symmetric frequency map also integrates in every degree when the characteristic exceeds the polynomial degree. Keeping the average squared Fourier modulus through each descent yields an explicit factorial bound.

Corollary II.247.11 (Polynomial phase approximation at every order). *Let $k \geq 2$, let $p \geq k$ be prime, and define*

$$K_k = 2 \prod_{r=1}^{k-2} (108r + 1) \leq 2 \cdot 109^{k-2} (k-2)!, \quad (\text{II.247.15})$$

with the empty product equal to one. Suppose $f : \mathbb{F}_p \rightarrow \mathbb{C}$ satisfies $|f| \leq 1$, $Q_k(f) \geq 1 - \varepsilon$, and $0 \leq \varepsilon \leq 1/K_k$. There exists $P \in \mathbb{F}_p[X]$ of degree at most $k - 1$ such that

$$|\mathbb{E}_x f(x) e_p(-P(x))| \geq \sqrt{1 - K_k \varepsilon} - \varepsilon. \quad (\text{II.247.16})$$

For a suitable unit scalar z ,

$$\|f - z e_p(P)\|_2^2 \leq 2(1 - \sqrt{1 - K_k \varepsilon}) + 2\varepsilon \leq (2K_k + 2)\varepsilon. \quad (\text{II.247.17})$$

Proof. Let u be unimodular with $Q_{r+2}(u) \geq 1 - D$, where $r \geq 1$. Lemma II.247.4 gives mean Fourier deficit at most D and additive-test failure at most $9D$. Lemma II.247.6 therefore supplies a scalar c for which the selected frequency agrees with $ch_1 \cdots h_r$ outside a set of proportion at most $108rD$. Using (II.247.6) on the agreement set gives

$$\mathbb{E}_{\mathbf{h}} \left| \widehat{\Delta_{h_1} \cdots \Delta_{h_r} u}(ch_1 \cdots h_r) \right|^2 \geq 1 - (108r + 1)D. \quad (\text{II.247.18})$$

If $p > r + 1$, put $u_1(x) = u(x) e_p(-cx^{r+1}/(r+1)!)$. The additive difference of order r of $cX^{r+1}/(r+1)!$ is

$$ch_1 \cdots h_r X + \gamma(\mathbf{h})$$

for some constant term $\gamma(\mathbf{h})$. Indeed each difference reduces the degree by one and multiplies the leading coefficient by the previous degree and its direction. The factorial is invertible under the stated characteristic assumption. The Gowers recursion and (II.247.18) now give the degree descent

$$Q_{r+1}(u_1) \geq 1 - (108r + 1)D. \quad (\text{II.247.19})$$

The symmetric-table lemma has no smallness premise, so no additional threshold is needed at any step.

Normalize f to its unit phase g . By Lemma II.247.1, $Q_k(g) \geq 1 - 2\varepsilon$. Apply (II.247.19) successively for $r = k - 2, k - 3, \dots, 1$. Every modulation has degree at most $k - 1 < p$, and their sum is a polynomial R of degree at most $k - 1$. The final function $v = g e_p(-R)$ satisfies

$$Q_2(v) \geq 1 - 2\varepsilon \prod_{r=1}^{k-2} (108r + 1) = 1 - K_k \varepsilon.$$

When $k = 2$ the sequence of modulations is empty and $R = 0$. Parseval supplies a frequency a such that $|\widehat{v}(a)|^2 \geq Q_2(v) \geq 1 - K_k \varepsilon$. Take $P = R + aX$ and return from g to f using $\|f - g\|_1 \leq \varepsilon$. This proves (II.247.16); scalar alignment and $1 - \sqrt{1 - x} \leq x$ give (II.247.17). Finally, $108r + 1 \leq 109r$ for $r \geq 1$ proves the upper bound in (II.247.15). $\square$

The generic fourth-order constant is $K_4 = 47306$. Corollary II.247.10 retains the sharper constant 6194 by using the two-variable collision argument and the direct one-variable BLR correction. The higher-characteristic condition above serves a specific purpose: the factorials used to integrate the frequency maps must be invertible. No analogous classical-polynomial conclusion is asserted here in the excluded characteristics.

Proposition II.247.12 (The order of the squared error cannot be improved). *Fix $k \geq 2$ and a prime $p \geq k$. A bound $o(\varepsilon)$ cannot uniformly replace the order $O(\varepsilon)$ of the squared L^2 distance in Corollary II.247.11, even on this fixed field. The same statement applies to Corollary II.247.10 when $k = 4$.*

Proof. Take a nonconstant real $h : \mathbb{F}_p \rightarrow \mathbb{R}$ of mean zero, put $\sigma^2 = \mathbb{E}h^2 > 0$, and let $f_t(x) = \exp(ith(x))$. There are finitely many polynomial phase functions of degree at most $k-1$ modulo a constant scalar. Every nonconstant one has mean of modulus strictly less than one. At $t=0$ the constant phase is strictly closer than all these alternatives; continuity preserves this for sufficiently small t . Hence

$$\inf_{|z|=1, \deg P \leq k-1} \|f_t - ze_p(P)\|_2^2 = 2 - 2|\mathbb{E}e^{ith}| = \sigma^2 t^2 + O(t^3).$$

In the k th-order cube product, write the additive alternating sum of the 2^k values of h as S . Taylor expansion gives $1 - Q_k(f_t) = (t^2/2)\mathbb{E}S^2 + O(t^3)$. Distinct cube vertices are pairwise independent uniform points: the corresponding two affine forms have rank two over $\mathbb{F}_p$. Their mixed products of mean-zero h average to zero, while each diagonal term contributes σ^2 . Thus $\mathbb{E}S^2 = 2^k \sigma^2$, and

$$1 - Q_k(f_t) = 2^{k-1} \sigma^2 t^2 + O(t^3).$$

The ratio of the optimal squared distance to this deficit tends to 2^{1-k} , in particular $1/8$ for $k=4$. This establishes the error order, without identifying an optimal global numerical constant. $\square$

II.248 Sources 59 and 56: affine rigidity on larger ranges and relative repair [write]

Added in the tenth write, 7 October 2026.

The next twelve sections are source 59's Sections 7–10 and source 56's Sections 2–9. *Logically they continue source 11's chapter* on affine rigidity at the Section 6 endpoint (Section II.22) and source 12's extension (Section II.33); they are printed here because no number printed earlier may move. The introductions are Sections F.64 and F.61; source 59's Sections 2–6 are the next chapter (Section II.261).

What source 59 proves here. For $f : G \rightarrow W$ at error density t from any affine reference, with $\varepsilon(f) = 1 - \mathcal{E}(f)/N^3$ and $\Psi(t) = 4t - 10t^2 + 6t^3$ (source 11's Ψ_0): an exact signed graph expansion (Lemma II.249.1) gives $\varepsilon \geq \Psi(t) + 6(t^3 - E(S)/N^3) + 2(1 - 3t)a$ for $t \leq 1/2$ when $W[2] = 0$ (Theorem II.249.3), improving source 11's colour coefficient $2(1 - 7t)$; averaging the signed kernel gives $\varepsilon \geq \Psi(t)$ on $t < 4/9$ ($W[2] = 0$) and on $t < 1/2$ ($W[2] = W[3] = 0$), with equality exactly for one nonzero error value on a subgroup coset (Theorem II.250.4; source 11 had $t < 1/7$), and both endpoints resolved (Corollary II.250.6). With the published correction theorem of Ben-Or, Coppersmith, Luby and Rubinfeld, as source 59 credits, it obtains recovery for $\varepsilon < 2/9$ (Theorem II.251.2; source 11: $1/9$), the exact non-affine energy $N^3 - 4N^2 + 10N - 6$ for $N \geq 16$ (Corollary II.251.3; source 11: $N \geq 34$), and the partial-graph threshold $\alpha^4 > 7/9$ at Proposition 6.1 (Corollary II.251.4; source 11: $8/9$); a map $\mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/3\mathbb{Z}$ of energy 1296 $>$ 1266 shows that the cutoff is at least 13 (Section II.252). Source 59 says that these global consequences follow from the prior cubic and classical correction and are not independently novel. Its Section II.252 also shows that a displayed probability estimate in the published proof of Green and Tao's Lemma 8.3 fails for an element of order five ($3/5$, not at most $1/2$): a counterexample to that estimate only, not to the lemma or to the inverse theorem, as source 59 says (note there).

What source 56 proves. For a weight $w : G \rightarrow [0, 1]$ of mean α on a finite abelian group with $u^4 = \|w - \alpha\|_{U^2}^4 \leq \alpha^5/1024$, any map to an arbitrary abelian group whose weighted additive-quadruple error is $\varepsilon \leq 1/1024$ agrees with a unique affine map off relative weighted mass $13\varepsilon/50$ (Theorem F.61.1); the domain condition does not depend on ε . A two-sided profile gives the coefficient

$1/4 + o(1)$, optimal (Theorem F.61.3, Corollary II.258.2); consequences are exact extension, a Fourier certificate and a four-query tester (Section II.259), and affine repair on quadratic level sets of rank $p^{R-4r} \geq 2^{14}$ (Theorem II.260.2). Linear repair is classical (Blum–Luby–Rubinfeld, Kazhdan–Ziegler), as source 56 says; its contribution is the explicit weighted relative theorem.

Relation to the formal statements (checked at HEAD `fe7165a3c`). `proposition_6_1`, `lemma_6_2` and `corollary_7_6` (Sections `06_07.lean`, lines 37, 47 and 197) and `lemma_10_9` (Section `10.lean`, line 335) have exact companions; source 56 says that it does not prove a stronger conclusion from the hypotheses of `lemma_10_9`. There is no formal counterpart of the affine-rigidity theory; nothing in this chapter is formalized.

II.249 Signed graph expansions and stronger affine rigidity

Source 59, delivered Section 7.

The following results refine the local graph-energy analysis in source 11 of [240], which quantifies the full-energy endpoint preceding Gowers’s Lemma 6.2 [1].

Let G be a finite abelian group of order N and let W be an abelian group with no nonzero element of order two. The target need not be finite: every sum used below is supported on the finite image of the map in question. For $f : G \rightarrow W$ write

$$\mathcal{E}(f) = \# \left\{ (x_1, x_2, x_3, x_4) \in G^4 : \begin{array}{l} x_1 + x_2 = x_3 + x_4, \\ f(x_1) + f(x_2) = f(x_3) + f(x_4) \end{array} \right\}.$$

The normalized graph defect is

$$\varepsilon(f) = 1 - \frac{\mathcal{E}(f)}{N^3}.$$

An affine map is a homomorphism plus a constant. Subtracting an affine map preserves $\mathcal{E}$. Fix such a map ℓ and set $g = f - \ell$, $S = \{x : g(x) \neq 0\}$, $s = |S|$, $t = s/N$. For $b \neq 0$ define

$$c_b = |g^{-1}(b)|, \quad M = \sum_b c_b^2, \quad A = s^2 - M, \quad C = \sum_b c_b c_{-b}, \quad Z = \sum_{a, b \neq 0, a+b \neq 0} c_a c_b c_{a+b}.$$

Here A counts ordered pairs of errors with different values, while C counts ordered opposite-value pairs. Set $a = A/N^2$, $c = C/N^2$ and $e_S = E(S)/N^3$, with $E(S)$ denoting ordinary set additive energy. The cubic profile is

$$\Psi(t) = 4t - 10t^2 + 6t^3.$$

II.249.1 An exact signed expansion

For a finitely supported real function u on an abelian group define

$$\mathcal{E}(u) = \sum_{z_1+z_2=z_3+z_4} u(z_1)u(z_2)u(z_3)u(z_4) = \sum_z (u * u)(z)^2.$$

In $G \times W$, let $B = 1_{G \times \{0\}}$ and

$$D(x, w) = 1_S(x)(1_{\{g(x)\}}(w) - 1_{\{0\}}(w)).$$

Thus $B + D$ is exactly the indicator of the graph of g .

Lemma II.249.1 (Signed graph identity). *For every such g ,*

$$\begin{aligned}\mathcal{E}(g) = & N^3 - 4sN^2 + N(6s^2 + 4M + 2C) \\ & - 4s^3 - 8sM - 4sC + 4Z + \mathcal{E}(D).\end{aligned}\tag{II.249.1}$$

Moreover,

$$\mathcal{E}(D) \leq 6E(S).\tag{II.249.2}$$

Proof. Expand the four factors in $\mathcal{E}(B + D)$. The term with no D is N^3 . For a term with one D , the target coordinate of that factor must be zero, and the other two free domain coordinates contribute N^2 ; its value is $-sN^2$, giving the linear term.

Put $m_0 = -s$ and $m_b = c_b$ for $b \neq 0$. With two D factors in opposite-sign positions, the contribution is $N \sum_b m_b^2 = N(s^2 + M)$. There are four choices. For a same-sign pair the contribution is $N \sum_b m_b m_{-b} = N(s^2 + C)$, and there are two choices. With three D factors, the fourth domain point is uniquely determined and the target condition gives

$$\sum_{a+b=c} m_a m_b m_c = Z - s^3 - 2sM - sC.$$

There are four identical contributions, by exchanging the two positive positions, the two negative positions, or the two sides. The four- D term is $\mathcal{E}(D)$. This proves the identity.

For the bound, fix a domain additive quadruple inside S and let the four nonzero signed target labels be $y_1 = g(x_1)$, $y_2 = g(x_2)$, $y_3 = -g(x_3)$, $y_4 = -g(x_4)$. Its contribution to $\mathcal{E}(D)$ is

$$K(y_1, y_2, y_3, y_4) = 1 + P - T + F,$$

where P is the number of zero-sum pairs among the four positions, T the number of zero-sum triples, and F is 1 when the sum of all four labels is zero and 0 otherwise. There are no singleton terms because every y_i is nonzero.

Since W has no nonzero involution, the graph joining opposite labels is a disjoint union of complete bipartite graphs. On four vertices it has at most four edges, so $P \leq 4$. Thus $K \leq 1 + 4 + 1 = 6$. There are $E(S)$ relevant domain quadruples. Summing the pointwise upper bound proves (II.249.2); individual kernel values need not be nonnegative. $\square$

Lemma II.249.2 (Weighted Schur bounds). *One has $0 \leq C \leq A$ and*

$$Z \leq \frac{sA}{2}.\tag{II.249.3}$$

If W also has no nonzero element of order three, then

$$Z \leq \frac{s(A - C)}{2}.\tag{II.249.4}$$

Proof. Every opposite-value pair has distinct values, proving $C \leq A$. For an ordered triple of nonzero labels (u, v, w) , consider the three equations $u + v = w$, $u + w = v$, $v + w = u$. At most one can hold: two of them would force one nonzero label to have order two. Let R be the number that hold and let U count unequal pairs among the three positions. If $R = 1$ then $U \geq 2$, since $u = v = w$ would force the common value to be zero. Thus $2R \leq U$ pointwise. Sum with weight $c_u c_v c_w$. The two sides sum to $6Z$ and $3sA$, respectively, which proves (II.249.3).

For the stronger assertion let V count pairs that are neither equal nor opposite. In a solution $u + v = w$, the pair (u, v) is never opposite. If $u = v$, then each of (u, w) and (v, w) is neither equal nor opposite, since $w = 2u$, $u \neq 0$, and $3u \neq 0$. If $u \neq v$, the first pair contributes to V . Neither of the remaining pairs can be equal; if both were opposite we would have $u = v = -w$, contrary to $u \neq v$. Hence $V \geq 2$. Thus $2R \leq V$, whose weighted sum is $6Z \leq 3s(A - C)$. $\square$

Theorem II.249.3 (A stronger local cubic and structural remainder). *If $0 \leq t \leq 1/2$, then*

$$\varepsilon(f) \geq \Psi(t) + 6(t^3 - e_S) + 2(1 - 3t)a. \quad (\text{II.249.5})$$

In particular, for $0 < t < 1/3$, $\varepsilon(f) \geq \Psi(t)$, with equality precisely when $g = b1_S$ for one nonzero value b and S is a subgroup coset. If W has no nonzero elements of orders two or three and $0 < t < 2/5$, the same cubic and the same equality classification hold. More precisely,

$$\varepsilon(f) \geq \Psi(t) + 6(t^3 - e_S) + (4 - 10t)a + (6t - 2)c. \quad (\text{II.249.6})$$

Proof. For $g_0 = b1_S$ with $b \neq 0$, either the same expansion or direct support-pattern counting gives

$$\mathcal{E}(g_0) = N^3 - 4sN^2 + 10s^2N - 12s^3 + 6E(S).$$

Subtract (II.249.1) and use (II.249.2) to obtain the useful intermediate bound

$$\mathcal{E}(g_0) - \mathcal{E}(g) \geq (N - 2s)(4A - 2C) - 4Z. \quad (\text{II.249.7})$$

Since $s \leq N/2$, the inequalities $C \leq A$ and $Z \leq sA/2$ imply that its right side is at least $2(N - 3s)A$. This proves (II.249.5).

When $s < N/3$, both remainder terms are nonnegative and the coefficient of A is positive. Equality forces $A = 0$ and $E(S) = s^3$. The first condition makes all error values equal. For completeness, the equality $E(S) = s^3$ says that every $x - y + z$, with $x, y, z \in S$, lies in S . Fixing y shows that $S - y$ is a finite subgroup, so S is a coset. The converse follows from the displayed exact formula.

With no order-three elements, use (II.249.4) in (II.249.7) without first replacing C :

$$(N - 2s)(4A - 2C) - 2s(A - C) = (4N - 10s)A + (6s - 2N)C.$$

This proves (II.249.6). If $s \leq N/3$, its color part is at least $(2N - 4s)A$, using $C \leq A$ and $6s - 2N \leq 0$. If $N/3 \leq s < 2N/5$, it is at least $(4N - 10s)A$, using $C \geq 0$. In either range the coefficient is strictly positive. Equality is therefore classified exactly as above. $\square$

Relative to source 11 of [240], the first inequality replaces the earlier color coefficient $2(1 - 7t)$ by $2(1 - 3t)$ and enlarges the corresponding sufficient support radius from $1/7$ to $1/3$. The absence of order-three elements is visible in a further explicit remainder, rather than only in a qualitative conclusion. These radii are sufficient radii; no optimality assertion is made about them. Equality of the cubic is realized at every available subgroup-coset density in the stated ranges.

Corollary II.249.4 (Quantitative near-equality below one third). *If $0 < t < 1/3$ and $\varepsilon(f) \leq \Psi(t) + \eta$, then*

$$t^3 - e_S \leq \frac{\eta}{6}, \quad a \leq \frac{\eta}{2(1 - 3t)}.$$

For some nonzero b , the proportion of points having nonzero error different from b is at most $\eta/[2t(1 - 3t)]$.

Proof. Drop the other nonnegative remainder in (II.249.5). If $m = \max_b c_b$, then $\sum_b c_b^2 \leq sm$, so $s - m \leq A/s$. Divide by N . $\square$

II.250 Averaging the signed kernel: radii four ninths and one half

Source 59, delivered Section 8.

The exact expansion also admits a second use of the quartic term. Averaging three choices of signs gives a cubic bound in the color multiplicities. Retaining the contribution from missing fourth points then preserves a support-energy remainder while enlarging the local range.

Retain all notation from Section II.249 and put

$$Z_3 = \sum_{b \neq 0, 3b=0} c_b^2 c_{-b}, \quad z_3 = Z_3/N^3.$$

Lemma II.250.1 (A torsion-resolved Schur estimate). *One has*

$$Z \leq \frac{s(A - C)}{2} + Z_3, \quad Z_3 \leq \frac{sC}{2}.$$

Proof. In the proof of Lemma II.249.2, the only way a Schur solution can fail to supply two pairs that are neither equal nor opposite is to have two labels equal to b and the third equal to $-b$, where $3b = 0$. For clarity, in $u + v = w$, if $u = v$ and $w = -u$ then $3u = 0$; if $u \neq v$, then (u, v) already contributes, and both other pairs cannot be opposite. Thus the pointwise inequality is $2R \leq V + 2J$, where J indicates such an exceptional triple. The weighted sums of R, V, J are $3Z, 3s(A - C)$ and $3Z_3$. This proves the first inequality.

Group the nonzero order-three labels into pairs $\{b, -b\}$. Such a pair contributes $c_b c_{-b}(c_b + c_{-b})$ to Z_3 and $2c_b c_{-b}$ to C . Since $c_b + c_{-b} \leq s$, summation proves the second inequality. $\square$

Lemma II.250.2 (Averaged signed-kernel bound). *For every finite error support and every target without nonzero involutions,*

$$\mathcal{E}(D) \leq 6s^3 - 2sA + sC - Z_3. \quad (\text{II.250.1})$$

Proof. Use the kernel $K = 1 + P - T + F$ from the proof of Lemma II.249.1. If $P = 4$, the four signed labels are two copies of b and two of $-b$ for some nonzero b . If $P = 3$, their opposite-pair graph must be the complete bipartite graph $K_{1,3}$; hence the labels are three copies of b and one of $-b$, and $F = 0$, since $2b \neq 0$. If $P \leq 2$, use $F \leq 1$. Consequently

$$K \leq 6, \quad K \leq 4 \text{ unless the four labels are balanced copies of } \{b, -b\}. \quad (\text{II.250.2})$$

If three specified signed labels are all equal to b and $3b = 0$, then $T \geq 1$. The balanced case is impossible; the preceding $P = 3$ and $P \leq 2$ alternatives therefore both give $K \leq 3$. If three specified signed labels are all equal but $3b \neq 0$, they still cannot be extended to a balanced quadruple, giving $K \leq 4$.

Fix an ordered triple $(x_1, x_2, x_3) \in S^3$, with nonzero error labels (u, v, w) . There are three ways to designate one of these three positions as negative and the other two as positive. In each case the fourth domain point, in the remaining negative position, is uniquely determined. If it is outside S , its contribution is zero. Otherwise apply the kernel caps above. The caps used below are positive, so they also bound the zero contribution when the fourth point is outside S .

For each designation start with the upper bound 6 and call the amount subtracted from this bound its guaranteed loss. Let U be the number of unequal pairs among (u, v, w) , let V count pairs that are neither equal nor opposite, and let J indicate that the triple uses both labels of one pair $\{b, -b\}$ with $3b = 0$. The sum of the three guaranteed losses is at least $U + V + J$:

- If $u = v = w$, then $U = V = J = 0$, and no loss is required.
- If the triple uses both labels of a single pair $\{b, -b\}$, exactly one of the three designations makes its three signed labels all equal. This designation loses at least 2, and loses at least 3 when $3b = 0$. Here $U = 2$, $V = 0$, so the total is at least $U + V + J$.
- If the triple meets at least two distinct pairs $\{b, -b\}$, no designation can extend to a balanced quadruple. Every designation loses at least 2, for a total loss at least 6. Since $U \leq 3$, $V \leq 3$ and $J = 0$, this is sufficient.

For each of the three designations separately, summing the actual contribution over all triples in S^3 equals $\mathcal{E}(D)$: this is just a reparametrization of the four-factor energy, with the three free positions assigned the indicated signs. Therefore averaging the three upper bounds gives

$$\mathcal{E}(D) \leq 6s^3 - \frac{1}{3} \sum_{x_1, x_2, x_3 \in S} (U + V + J).$$

The three sums are respectively $3sA$, $3s(A - C)$ and $3Z_3$. This proves (II.250.1). $\square$

Lemma II.250.3 (The support deficit in the averaged bound). *Under the hypotheses of Lemma II.250.2,*

$$\mathcal{E}(D) \leq 6s^3 - 2sA + sC - Z_3 - 3(s^3 - E(S)). \quad (\text{II.250.3})$$

If W also has no nonzero order-three elements, then

$$\mathcal{E}(D) \leq 6s^3 - 2sA + sC - 4(s^3 - E(S)). \quad (\text{II.250.4})$$

The coefficients 3 and 4 cannot be increased in these respective inequalities while keeping the other terms fixed.

Proof. In the preceding proof let L_i be the cap assigned to designation $i \in \{1, 2, 3\}$ of an ordered triple, and let K_i be its actual contribution. Thus $L_i \in \{3, 4, 6\}$ and

$$\sum_{i=1}^3 L_i \leq 18 - (U + V + J).$$

Write $I_i = 1$ when the fourth point for designation i is outside S , and $I_i = 0$ otherwise. Since $K_i = 0$ when $I_i = 1$ and every cap is at least 3, one has $K_i \leq L_i - 3I_i$. For each designation separately,

$$\sum_{(x_1, x_2, x_3) \in S^3} I_i = s^3 - E(S).$$

Summing the three inequalities and dividing by three proves (II.250.3). In the absence of order three, the cap 3 never occurs and $J = Z_3 = 0$. Replacing $3I_i$ by $4I_i$ proves (II.250.4).

For sharpness, take $G = \mathbb{Z}/5\mathbb{Z}$, $S = \{0, 1\}$, $g(0) = b \neq 0$, $g(1) = -b$, and $g = 0$ elsewhere. Directly,

$$s = 2, \quad E(S) = 6, \quad A = C = 2, \quad \mathcal{E}(D) = 36.$$

For a target element b of order three, $Z_3 = 2$ and the original kernel bound has right side 42; for a target without order three it has right side 44. Since $s^3 - E(S) = 2$, both stated support coefficients attain equality in the corresponding examples. $\square$

Theorem II.250.4 (Larger local ranges with color and support remainders). *For every target without nonzero involutions,*

$$\varepsilon(f) - \Psi(t) \geq 3(t^3 - e_S) + (4 - 8t)a + (5t - 2)c - 3z_3. \quad (\text{II.250.5})$$

Consequently, if $0 < t < 4/9$ then

$$\varepsilon(f) \geq \Psi(t) + 3(t^3 - e_S) + \left(2 - \frac{9t}{2}\right)a \geq \Psi(t). \quad (\text{II.250.6})$$

If W also has no nonzero order-three elements, the sufficient range improves to $0 < t < 1/2$, with

$$\varepsilon(f) \geq \Psi(t) + 4(t^3 - e_S) + \min\{2 - 3t, 4 - 8t\}a. \quad (\text{II.250.7})$$

In either range, equality with $\Psi(t)$ occurs precisely for one nonzero error value supported on a subgroup coset.

Proof. The signed identity rearranges exactly as

$$N^3(\varepsilon(f) - \Psi(t)) = (N - 2s)(4A - 2C) - 4Z + 6s^3 - \mathcal{E}(D).$$

Apply Lemmas II.250.1 and II.250.3. Retaining the term $3(s^3 - E(S))$, the remaining right side is at least

$$(N - 2s)(4A - 2C) - 2s(A - C) - 4Z_3 + 2sA - sC + Z_3 = (4N - 8s)A + (5s - 2N)C - 3Z_3,$$

proving (II.250.5). Since $Z_3 \leq sC/2$, this is at least

$$(4N - 8s)A + \left(\frac{7s}{2} - 2N\right)C.$$

When $s < 4N/9$, the coefficient of C is negative, so $C \leq A$ gives $(2N - 9s/2)A$, proving (II.250.6).

In the no-order-three case, use the stronger support remainder $4(s^3 - E(S))$. The coefficient $5s - 2N$ changes sign at $s = 2N/5$. Below that point use $C \leq A$, giving $(2N - 3s)A$; above it use $C \geq 0$, giving $(4N - 8s)A$. Both coefficients are positive in the stated range, and they agree at $s = 2N/5$. This proves (II.250.7).

Equality with the cubic forces $A = 0$, hence $g = b1_S$ with $b \neq 0$. The exact one-valued-error formula then gives $\varepsilon(f) - \Psi(t) = 6(t^3 - e_S)$, so $E(S) = s^3$ and S is a subgroup coset. Conversely, these configurations attain equality. $\square$

[write] **7 October 2026; sources 54–59.** *Against source 11* (Section II.22). Source 11's Theorem F.16.1 proves $\varepsilon \geq \Psi_0(t)$ with coset equality on $0 < t < 1/7$ for targets without involutions (and Ψ_1 on $t < 1/3$ with them); this theorem extends the first range to $t < 4/9$, and to $t < 1/2$ without 3-torsion, with the same equality cases, which answers the first part of source 11's question how far the radius $1/7$ extends (item II.19); the involution case is this source's Question V.223.8.

[write] **7 October 2026; sources 60–69.** *The whole range for elementary abelian 3-group targets* (source 67, Section II.293): for these targets, which have no involutions but have 3-torsion, $\varepsilon \geq \Psi(t)$ holds for every $t \leq 2/3$, with the same coset equality cases for $0 < t < 2/3$ and a separate affine family at $t = 2/3$ (its Theorem II.294.1). Source 67 credits this theorem's ranges $4/9$ and $1/2$ to source 59.

Lemma II.250.5 (Rigidity of the averaged kernel bound). *Assume $s > 0$ and $W[2] = \{0\}$. Equality in (II.250.1) holds if and only if $g = b1_S$ for a nonzero $b \in W$ and S is a subgroup coset.*

Proof. The proof of Lemma II.250.2 assigns each of the three designations an upper cap belonging to $\{3, 4, 6\}$. Its final deficit is a sum of nonnegative pointwise deficits: the actual contribution is at most its assigned cap, and the sum of the three caps is at most $18 - (U + V + J)$. Equality of the total therefore forces equality at every one of these finite comparisons.

Every cap is positive. Thus no fourth point can be outside S , since its actual contribution would then be zero. It follows that S is closed under $x + y - z$, and hence is a coset. Equivalently, this also follows from Lemma II.250.3.

If two occurring colors u, v are neither equal nor opposite, use a triple of points with colors (u, u, v) , allowing a repeated point for the two copies of u . Every designation has cap 4, so the sum of caps is 12. But $U = V = 2$ and $J = 0$, giving $18 - (U + V + J) = 14$. This is strict slack, a contradiction. All colors therefore belong to one pair $\{b, -b\}$.

For each occurring color b , let $T_b = \{x : g(x) = b\}$. On a triple from T_b^3 , all three caps are 6. Equality forces the actual kernel to be 6, which requires a balanced quadruple. Its fourth point therefore also has color b . Consequently T_b is closed under $x + y - z$, so it is a coset.

Suppose both b and $-b$ occur. Choose $x \in T_b$ and $z \in T_{-b}$. Since S is a coset, $2x - z \in S$. Since T_b is a coset and $z \notin T_b$, one has $2x - z \notin T_b$, so $g(2x - z) = -b$. In the triple (x, x, z) designate z as negative. All four signed labels are then b . There are no zero-sum pairs, and the four-label sum is nonzero: $4b = 0$ would contradict $W[2] = \{0\}$. The actual kernel is consequently

$$1 - 4\mathbf{1}_{3b=0},$$

which is 1 or -3 . Its assigned cap is respectively 4 or 3, again a strict deficit. Thus only one color occurs.

Conversely, if $g = b\mathbf{1}_S$ with S a coset, then $A = C = Z_3 = 0$ and $\mathcal{E}(D) = 6E(S) = 6s^3$, proving equality. $\square$

Corollary II.250.6 (The two local endpoints). *If $W[2] = \{0\}$ and $t = 4/9$, then*

$$\varepsilon(f) > \Psi(4/9).$$

If $W[2] = W[3] = \{0\}$ and $t = 1/2$, then $\varepsilon(f) \geq \Psi(1/2) = 1/4$, with equality precisely when $g = b\mathbf{1}_S$ for a nonzero b and an index-two subgroup coset S .

Proof. The comparisons proving Theorem II.250.4 remain valid at their closed endpoints. At $t = 4/9$ they give

$$\varepsilon(f) - \Psi(t) \geq 3(t^3 - e_S).$$

If equality with $\Psi(t)$ held, then $E(S) = s^3$ and S would be a coset. A coset in a finite group has reciprocal-integer density, which cannot equal $4/9$. The inequality is therefore strict.

Now assume $W[3] = \{0\}$ and $N = 2s$. The exact signed identity can be written as a sum of nonnegative terms:

$$\begin{aligned} N^3(\varepsilon(f) - \Psi(t)) &= 4 \left(\frac{s(A - C)}{2} - Z \right) \\ &\quad + (6s^3 - 2sA + sC - \mathcal{E}(D)) + sC. \end{aligned}$$

The Schur bound and (II.250.1) prove nonnegativity. Equality forces equality in the latter kernel bound. Lemma II.250.5 therefore gives one nonzero color on a coset, whose density $1/2$ makes its index two. Conversely, the exact one-valued formula gives equality for every such map. $\square$

Corollary II.250.7 (Near-equality on the enlarged odd-target range). *If $0 < t < 4/9$, $\eta \geq 0$, and $\varepsilon(f) \leq \Psi(t) + \eta$, then*

$$a \leq \frac{2\eta}{4-9t}, \quad t^3 - e_S \leq \min \left\{ \frac{\eta}{3}, \frac{\eta}{6} + \frac{2(3t-1)_+\eta}{3(4-9t)} \right\}.$$

For some nonzero b , the fraction of all domain points with nonzero error different from b is at most $2\eta/[t(4-9t)]$. If $W[3] = \{0\}$, $0 < t < 1/2$, and $\varepsilon(f) \leq \Psi(t) + \eta$ with $\eta \geq 0$, then writing $h(t) = \min\{2-3t, 4-8t\} > 0$, one has

$$a \leq \frac{\eta}{h(t)}, \quad t^3 - e_S \leq \frac{\eta}{4},$$

$$\frac{1}{N} \min_{b \neq 0} |\{x : g(x) \notin \{0, b\}\}| \leq \frac{\eta}{th(t)}.$$

Proof. The first bound and $t^3 - e_S \leq \eta/3$ follow from (II.250.6). Since $t < 4/9 < 1/2$, (II.249.5) also applies, giving $6(t^3 - e_S) \leq \eta + 2(3t-1)_+a$. Substitute the first bound. The color conclusion again follows from $s - \max_b c_b \leq A/s$. The no-order-three conclusions follow in the same way from (II.250.7). $\square$

The proof identifies the role of order three explicitly. Such torsion does not change the original one-valued cubic, but it creates exceptional Schur triples and lowers a signed-kernel cap from four to three. The combined bookkeeping enlarges the sufficient local radius to $4/9$; removing this torsion enlarges it further to $1/2$. The support deficits have coefficients 3 and 4 in the corresponding kernel bounds, and Corollary II.250.6 settles equality at both endpoints. Neither number is asserted to be the best possible universal radius.

II.251 Classical correction and global affine recovery

Source 59, delivered Section 9.

The next lemma is the classical sharp-range majority correction argument, with its collision-to-mode estimate retained [27, 28]. It is included to give a complete proof of the global consequences for the repository's Section 6 interface. Neither majority decoding nor the $7/9$ homomorphism-test threshold is claimed as a new result.

Priority of the global consequences. The sharper published correction theorem of Ben-Or, Coppersmith, Luby, and Rubinfeld [27], restated as Theorem 4.2 by Diaconis, Shao, and Soundararajan [28], gives

$$\Pr[F \neq L] \leq \frac{3 - \sqrt{9 - 24\rho}}{12} < \frac{1}{7} \quad (\rho < 2/9).$$

Consequently, the global graph-defect range $\varepsilon < 2/9$ and the exact non-affine energy formula for $N \geq 16$ below already follow from that published theorem and source 11's earlier local cubic on $t < 1/7$ [240]. They are not independently novel global results. Their role here is to complete and simplify an explicitly unfinished correction/cutoff step in the pinned report. The new local color and support estimates of the preceding sections are separate refinements. We reproduce an elementary correction proof with a weaker intermediate distance bound, which is sufficient for the complete argument below.

Lemma II.251.1 (Majority correction in the 7/9 range). *Let $F : G \rightarrow W$ and suppose that*

$$\rho = \Pr_{x,y}[F(x+y) \neq F(x) + F(y)] < 2/9.$$

Put $q = (1 - \sqrt{1 - 4\rho})/2$. There is a homomorphism $L : G \rightarrow W$ with $\Pr[F \neq L] \leq q < 1/3$.

Proof. For fixed h , let $Z_h(y) = F(h+y) - F(y)$. For independent uniform y, z , both

$$F(h+y) + F(z) = F(h+y+z), \quad F(y) + F(h+z) = F(h+y+z)$$

hold except on a set of probability at most 2ρ . Each tested pair is uniform on G^2 . Thus the collision probability of Z_h is at least $1 - 2\rho > 1/2$.

For any probability distribution with largest atom p , one has $\sum p_i^2 \leq p$, so this largest atom exceeds $1/2$. Also $\sum p_i^2 \leq p^2 + (1-p)^2$, because the remaining atoms have total mass $1-p$. Therefore $p(1-p) \leq \rho$, and the branch $p > 1/2$ gives $p \geq 1-q$. Let $L(h)$ be this unique modal value.

Since $q < 1/3$, the three identities

$$Z_{h+k}(x) = L(h+k), \quad Z_h(x+k) = L(h), \quad Z_k(x) = L(k)$$

hold simultaneously for some x . Their exact cocycle identity proves $L(h+k) = L(h) + L(k)$. If $F(h) \neq L(h)$, at least a fraction $1-q$ of the values $Z_h(y)$ differ from $F(h)$. Averaging yields $\rho \geq (1-q)\Pr[F \neq L]$. Since $\rho = q(1-q)$, the conclusion follows, including the case $\rho = 0$. $\square$

Let $D(u)$ be the inverse of Ψ on $[0, 1/7]$; this restriction is strictly increasing. In the next statement $d(f)$ denotes the smallest normalized Hamming distance from f to an affine map.

Theorem II.251.2 (Global sharp-profile recovery on a larger range). *If $\varepsilon(f) < 2/9$, then the nearest affine map is unique and*

$$d(f) \leq D(\varepsilon(f)).$$

For positive defect, equality holds precisely when the error from that map is one nonzero constant on a subgroup coset.

Proof. The average, over a, x, y , of the acceptance condition

$$f(a+x) + f(a+y) = f(a) + f(a+x+y)$$

is $1 - \varepsilon(f)$: the parametrization by (a, x, y) is a bijection onto ordered additive quadruples. Hence for some anchor a the map $F_a(x) = f(a+x) - f(a)$ has additive-test failure probability $\rho \leq \varepsilon(f)$. Lemma II.251.1 produces an affine map within distance less than $1/3$ from f .

Thus $t = d(f) < 1/3$, and Theorem II.249.3 gives $\varepsilon(f) \geq \Psi(t)$. The derivative $\Psi'(t) = 4 - 20t + 18t^2$ has just one zero in $[1/7, 1/3]$, and it is a maximum of Ψ . Consequently

$$\min_{[1/7, 1/3]} \Psi = \min\{132/343, 4/9\} > 2/9.$$

Therefore $t < 1/7$, where inversion gives the stated estimate.

Two distinct affine maps into W agree either nowhere or on a coset of the kernel of a nonzero homomorphism. The image of that homomorphism is a nontrivial finite subgroup of W , hence has order at least three. Their mutual distance is therefore at least $2/3$. Two nearest maps within distance less than $1/3$ from f must coincide. Finally, equality in the inverse estimate is equality in the local cubic, whose equality case has already been proved. The converse follows from the exact one-valued-error formula. $\square$

Corollary II.251.3 (Exact non-affine energy for $N \geq 16$). *Suppose that $W \neq \{0\}$ and $N \geq 16$. Then*

$$\max_{f \text{ non-affine}} \mathcal{E}(f) = N^3 - 4N^2 + 10N - 6.$$

The maximizers are precisely the one-point modifications of affine maps.

Proof. For $N \geq 16$, monotonicity in N gives

$$\Psi(1/N) \leq \Psi(1/16) = \frac{435}{2048} < \frac{2}{9}.$$

If a non-affine map had energy greater than the claimed value, the preceding theorem would give $d(f) < 1/N$, impossible on an N -point domain. At equality it gives $d(f) \leq 1/N$, hence $d(f) = 1/N$. Conversely, an affine map changed at one point by a nonzero value has the displayed energy, by the exact one-valued-error formula. It is non-affine because distinct affine maps cannot differ at a single point when $N \geq 16$. $\square$

[write] **7 October 2026; sources 54–59.** *Checked at the write* (exact rationals). $\Psi(1/16) = 435/2048 < 2/9$; $\Psi(1/7) = 132/343$ and $\Psi(1/3) = 4/9$, both above $2/9$; $\Psi'(t) = 4 - 20t + 18t^2 > 0$ on $[0, 1/16]$; $\Psi(1/2) = 1/4$, so $(1 - \Psi(1/2)) \cdot 12^3 = 1296$ against $12^3 - 4 \cdot 12^2 + 10 \cdot 12 - 6 = 1266$ (Section II.252). *Against source 11's cutoff* (Theorem F.16.3, $N \geq 34$): for targets without involutions the cutoff is now 16, and the third write's exhaustive table (Remark W2: $\mathbb{Z}/6 \rightarrow \mathbb{Z}/3$ and $\mathbb{Z}/9 \rightarrow \mathbb{Z}/3$ fail) together with this source's $\mathbb{Z}/12 \rightarrow \mathbb{Z}/3$ example puts the exact universal cutoff between 13 and 16; the intake's marked derivation of the cutoff 30 (item II.21) is superseded. This source credits the global consequence to classical correction and source 11's cubic. Source 67, placed later and not yet written in, records the maximum from $N \geq 13$ with 13 optimal for elementary abelian 3-group targets.

[write] **7 October 2026; sources 60–69.** For elementary abelian 3-group targets source 67 proves the formula from $N \geq 13$, with 13 optimal (its Corollary II.295.5); its sharpness example at $N = 12$ is this source's parity map $\mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/3\mathbb{Z}$ (energy $1296 > 1266$, Section II.252), which source 67 does not attribute to this source at that point, although it credits this source's ranges, recovery threshold $2/9$ and cutoff 16 (its bibliography entry [271]).

This improves the pinned repository's sufficient cutoff 34 to 16 for targets without involutions [240], as a consequence of classical correction and the prior local cubic. It does not claim that 16 is the exact cutoff. Classifying the exact maximum and all extremizers for domains of size $N < 16$ remains a separate problem.

Corollary II.251.4 (Partial graphs and the Section 6 interface). *Let $B \subseteq G$, let $\phi : B \rightarrow W$, and suppose that its graph Γ_B satisfies $E(\Gamma_B) \geq (1 - \varepsilon)N^3$ with $0 \leq \varepsilon < 2/9$. There is a unique affine map ℓ with at most $D(\varepsilon)N$ disagreements with ϕ on B . Moreover,*

$$\frac{|G \setminus B|}{N} \leq 1 - (1 - \varepsilon)^{1/3}.$$

Every arbitrary completion of ϕ gives the same recovered affine map.

Proof. Every completion has graph energy at least $E(\Gamma_B)$, and so the global theorem applies. The set bound $E(\Gamma_B) \leq |B|^3$ gives $\beta = |B|/N \geq (1 - \varepsilon)^{1/3}$. If two affine maps satisfy the disagreement bound, they agree at least on $(\beta - 2D(\varepsilon))N$ points. Since $\beta > 7/9$ and $D(\varepsilon) < 1/7$, this exceeds $(7/9 - 2/7)N = 31N/63 > N/3$, impossible for two distinct affine maps into a target without involutions. This also proves independence from the completion. $\square$

[write] **7 October 2026; sources 54–59.** Source 11’s Corollary II.32.2 enters the high-correlation regime under $\alpha^4 > 8/9$; this corollary lowers it to $7/9$, as the next paragraph says.

In the prime-cyclic frequency-selection setting of Gowers’s Proposition 6.1 [1], its graph-energy conclusion $E(\Gamma_B) \geq \alpha^4 N^3$ therefore enters this regime as soon as $\alpha^4 > 7/9$, replacing the earlier sufficient condition $\alpha^4 > 8/9$. The error is at most $D(1 - \alpha^4)N$. This strengthens a high-correlation endpoint consequence of the proposition; it changes neither its exponent four nor the general quantitative bound in Szemerédi’s theorem.

II.252 Examples, comparisons, and precise claim boundaries

Source 59, delivered Section 10.

II.252.1 A finite-field example where image size matters

Let $H = \mathbb{F}_2^3$ and let F consist of six of its seven nonzero elements. Here $q = 6$, so the exact universal rank is two. Two suitable linear functionals have common kernel equal to the line through the omitted element. Their common fibers each have size two. The finite-image argument therefore costs only $2^2 = 4$, giving a guaranteed retained weight of $w(B)/4$ for any graph with that defect set and any k . Charging continuous arc lengths instead would give only $w(B)/(5k)^2$. The example separates the role of rank from the role of character-image size.

For a character of image order 81, the prime-primary radius for $p = 3$ is $1/6$. At order $k = 8$, the exact single-coordinate factor is

$$\frac{\lceil 81/48 \rceil}{81} = \frac{2}{81},$$

which improves the continuous factor $1/48$. These finite corrections are substantial when the image is small, and can still help when it is larger than the localization denominator.

II.252.2 An exact small-domain obstruction

Take $G = \mathbb{Z}/12\mathbb{Z}$, $W = \mathbb{Z}/3\mathbb{Z}$, and let f equal one on odd residues and zero on even residues. Relative to the zero affine reference, its error support is an index-two coset and $t = 1/2$. The one-valued formula in Section II.249 gives

$$\mathcal{E}(f) = (1 - \Psi(1/2)) 12^3 = \frac{3}{4} 12^3 = 1296.$$

The one-point formula would give $12^3 - 4 \cdot 12^2 + 10 \cdot 12 - 6 = 1266$. The map is non-affine: its image has two elements, whereas an affine map into $\mathbb{Z}/3\mathbb{Z}$ has image size one or three. Thus a size restriction in Corollary II.251.3 cannot simply be removed, even without target involutions. The optimal universal cutoff is at least 13 and at most 16; the orders between them require further analysis.

II.252.3 The published order-five probability issue

Green–Tao’s published separation lemma, Lemma 8.3 on page 106 of [4], seeks a small set of characters whose radius- $1/4$ Bohr neighborhood meets a given set only at zero. Its proof states, for

a uniformly chosen character and a fixed nonzero element f , the estimate

$$\mathbb{P}(\|\chi(f)\|_{\mathbb{T}} < 1/4) \leq 1/2. \quad (\text{II.252.1})$$

For an element of order five, evaluation is uniform on $\{0, 1/5, 2/5, 3/5, 4/5\}$. The arc contains $0, 1/5, 4/5$, so the probability is $3/5$. Hence (II.252.1), as a statement for every nonzero element in every finite abelian group, is false. The same displayed step occurs in the inspected arXiv version.

The conclusion is deliberately limited: this is a counterexample to that probability estimate. It is not a counterexample to the lemma's conclusion with its original rank allowance, and it is not a disproof of the inverse theorem. The present article does not decide whether the original separation conclusion follows by a different argument. Lemma II.262.2 gives the correct order-sensitive distribution, and radius $1/5$ supplies a valid universal half-survival estimate. Combined with strict adaptive deletion, it yields the sharp statement actually proved here. No claim is made to be the first to have noticed the issue.

[write] **7 October 2026; sources 54–59.** *Recorded as stated* (standing rule of 4 October 2026). The arithmetic is checked at the write: for f of order five, $\chi(f)$ is uniform on $\{0, 1/5, 2/5, 3/5, 4/5\}$ and exactly $0, 1/5$ and $4/5$ lie within $1/4$ of an integer, so the probability is $3/5 > 1/2$; in the exact small-ball law of Lemma II.262.2, $\beta_5(1/4) = (2\lceil 5/4 \rceil - 1)/5 = 3/5$. This is a counterexample to the displayed estimate (II.252.1) read as a statement for every nonzero element of every finite abelian group, not to Lemma 8.3 of Green and Tao or to their inverse theorem, as this source says; the write did not re-read the published proof, and [4] is the report's entry for that paper.

II.252.4 Attribution and the status of each contribution

The following table is intended to prevent an integration from turning an explicit local refinement into an unsupported priority claim.

Topic	Established input	Contribution or consequence here
Character localization	Green–Tao separation and Bohr extraction, Lemmas 8.3 and 9.2.	Exact distributions, strict integer ranks, sharp universal/odd rank–radius pairs, order profiles.
Finite-field extraction	Gowers–Milićević random-kernel argument; source 24.	Adaptive/projective rank endpoints and their inclusion in the general finite-image formula.
Weights	Averaging is classical; source 31 already has weighted prime-cyclic extraction.	Weights survive the full arbitrary-target localization and exact image-size correction. Weighting itself is not claimed new.
Mixed torsion	Source 10 has a torsion budget; source 24 asks for an extension beyond finite fields.	An explicit target-independent polynomial bound, with optional dependence on actual defect orders.
Defect realizability	A finite relation presentation is the elementary construction mechanism.	Every symmetric set is realized at a prescribed order, with exact maximum cardinality, weight, and graph difference size.

Topic	Established input	Contribution or consequence here
Affine cubic	Source 11 already has Ψ , coset equality, and a radius $1/7$ for targets without involutions.	Stronger support/color remainder; sharp kernel support coefficients; radii $4/9$ and $1/2$, including endpoint rigidity.
Global correction	Ben-Or–Coppersmith–Luby–Rubinfeld; Diaconis–Shao–Soundararajan.	A complete derivation of the larger repository recovery range, cutoff 16 , and partial-graph consequence. These follow from the old cubic plus classical correction.
Energy to doubling	Reiher–Schoen Theorem 1.2.	Its explicit composition with arbitrary-target restriction; no new BSG theorem.
One-character recurrence	The lonely runner literature; Bedert’s theorem and Riesz-product argument.	Transfer to graph defects and an explicit ℓ -dissociated formula over abelian groups; no claimed improvement to Bedert’s general bound.

The strongest mathematical additions are the exact separation refinement, the signed kernel inequalities, and the consequences proved from them. Whether any of these formulations has an earlier equivalent elsewhere is a matter for a fuller specialist priority check. Their written proofs and stated comparisons can be assessed independently of that question.

II.253 Normalization and elementary Fourier facts

Source 56, delivered Section 2.

Let $\widehat{G}$ be the character group of G . The Fourier transform and convolution are normalized by

$$\widehat{f}(\chi) = \mathbb{E}_x f(x) \overline{\chi(x)}, \quad (f * g)(x) = \mathbb{E}_y f(y) g(x - y).$$

Character orthogonality gives inversion and Parseval:

$$f(x) = \sum_{\chi \in \widehat{G}} \widehat{f}(\chi) \chi(x), \quad \sum_{\chi \in \widehat{G}} |\widehat{f}(\chi)|^2 = \mathbb{E}_x |f(x)|^2.$$

These formulas use an ordinary sum, not an average, on the dual group. The U^2 norm is

$$\|f\|_{U^2(G)}^4 = \mathbb{E}_{x,h,k} f(x) \overline{f(x+h)} \overline{f(x+k)} f(x+h+k) = \sum_{\chi} |\widehat{f}(\chi)|^4.$$

The identity follows by inserting Fourier inversion and averaging the three independent variables. Each nonzero surviving term has all four character indices equal.

Lemma II.253.1 (Energy and centering). *For a real nonnegative w of mean α , one has*

$$\mathcal{Q}(w) = \sum_{\chi} |\widehat{w}(\chi)|^4 = \alpha^4 + \|w - \alpha\|_{U^2(G)}^4. \quad (\text{II.253.1})$$

If $0 \leq v \leq 1$ and $b = \mathbb{E}v$, then

$$\mathcal{Q}(v) \leq b^3, \quad \sum_{\chi \neq 1} |\widehat{v}(\chi)|^4 \leq b^3 - b^4 \leq b^3. \quad (\text{II.253.2})$$

Proof. The Fourier expansion of (F.61.1) gives the first equality; removing the trivial character gives the second. For (II.253.2), bound the fourth factor $v(a + b - c)$ in the defining average by one. The remaining three variables are independent. Subtracting the trivial Fourier term gives the centered bound. $\square$

For later use write

$$\mu = \mathbb{E}w^2, \quad \sigma = \max_{\chi \neq 1} |\widehat{w}(\chi)|, \quad (\text{II.253.3})$$

with $\sigma = 0$ if G is the trivial group. Since $0 \leq w \leq 1$,

$$\alpha^2 \leq \mu \leq \alpha, \quad \sum_{\chi \neq 1} |\widehat{w}(\chi)|^2 = \mu - \alpha^2, \quad \sigma \leq u.$$

The last inequality is the elementary comparison between the largest term and the fourth-power sum of nonnegative terms.

II.253.1 Why the weight bound matters

The restriction $w \leq 1$ is used in the energy cap and in bounding μ by α . A general bounded nonnegative weight can first be divided by its supremum; the test error and relative distance are invariant under this rescaling, but the stated uniformity hypothesis must then be checked for the rescaled weight. It would be incorrect to keep the original α^5 threshold unchanged under arbitrary rescaling.

II.253.2 A collision-to-mode inequality

Lemma II.253.2. *If (p_j) is a probability distribution with finite support and $m = \max_j p_j$, then*

$$1 - m \leq 1 - \sum_j p_j^2.$$

Proof. The inequality $\sum_j p_j^2 \leq m \sum_j p_j = m$ is immediate. $\square$

All modes in the proof have finite support, even when H is infinite. No measure-theoretic selection principle and no ordering of H is required for the mathematical existence statements. An implementation may use any fixed tie-breaking rule.

II.254 Autocorrelation, triangle variance, and pointwise links

Source 56, delivered Section 3.

The domain estimates are independent of the labels. Their role is to replace successive exceptional-set deletions by direct moment bounds. Define

$$r(h) = \mathbb{E}_x w(x)w(x+h), \quad (\text{II.254.1})$$

$$T(h, k) = \mathbb{E}_x w(x)w(x+h)w(x+h+k), \quad (\text{II.254.2})$$

$$D_w(x) = \mathbb{E}_{y,z} w(y)w(z)w(y+z-x). \quad (\text{II.254.3})$$

The symbol $D_w(x)$ is a domain link density, whereas $\mathcal{D}_w(\phi)$ is the label defect; they are different quantities.

Lemma II.254.1 (Autocorrelation identities). *One has*

$$0 \leq r(h) \leq \alpha, \quad \mathbb{E}_h r(h) = \alpha^2, \quad \mathbb{E}_h (r(h) - \alpha^2)^2 = u^4. \quad (\text{II.254.4})$$

Proof. Positivity is immediate and $w(x+h) \leq 1$ gives $r(h) \leq \alpha$. The mean follows by averaging x and $x+h$ independently. The square average $\mathbb{E}_h r(h)^2$ is $\mathcal{Q}(w)$, by the change of variables that identifies two pairs of the same difference with an additive quadruple. Therefore its variance is $\mathcal{Q}(w) - \alpha^4 = u^4$. $\square$

Lemma II.254.2 (Exact triangle second moment). *The triangle density satisfies*

$$\mathbb{E}_{h,k} T(h, k) = \alpha^3, \quad (\text{II.254.5})$$

$$\mathbb{E}_{h,k} T(h, k)^2 = \mathbb{E}_z r(z)^3. \quad (\text{II.254.6})$$

Consequently, for $V = \mathbb{E}_{h,k} (T(h, k) - \alpha^3)^2$,

$$V \leq \alpha(1+2\alpha)u^4, \quad \frac{\sqrt{V}}{\alpha^3} \leq \lambda\sqrt{1+2\alpha}. \quad (\text{II.254.7})$$

Proof. The variables $(x, x+h, x+h+k)$ are independently uniform, proving (II.254.5). For the second moment introduce two base points x, y and put $z = y - x$. The change of variables

$$a = x, \quad b = x+h, \quad c = x+h+k$$

is a bijection of G^3 . Thus

$$\begin{aligned} \mathbb{E}_{h,k} T(h, k)^2 &= \mathbb{E}_{z,a,b,c} w(a)w(a+z)w(b)w(b+z)w(c)w(c+z) \\ &= \mathbb{E}_z r(z)^3. \end{aligned}$$

Write $d(z) = r(z) - \alpha^2$. Then $\mathbb{E}d = 0$ and $\mathbb{E}d^2 = u^4$, so

$$V = 3\alpha^2 u^4 + \mathbb{E}d^3.$$

Since $d \leq \alpha - \alpha^2$, the pointwise inequality $d^3 \leq (\alpha - \alpha^2)d^2$ holds also when d is negative. This gives (II.254.7). Dividing by α^6 and using $u^4 = \lambda^2 \alpha^5$ proves the normalized estimate. $\square$

A slightly stronger bound replaces $\alpha(1+2\alpha)$ by $\mu + 2\alpha^2$, because Cauchy–Schwarz gives $r(h) \leq \mu$. We retain the simpler density-only form in the main constants.

Lemma II.254.3 (Uniform pointwise link bound). *For every $x \in G$,*

$$|D_w(x) - \alpha^3| \leq u^2 \sqrt{\mu - \alpha^2} \leq \lambda \alpha^3 \sqrt{1 - \alpha}. \quad (\text{II.254.8})$$

Proof. Set $\tilde{w}(x) = w(-x)$. Then $D_w = w * w * \tilde{w}$, whose Fourier coefficient at χ is $\widehat{w}(\chi)^2 \overline{\widehat{w}(\chi)}$. Its trivial coefficient is α^3 . Hence Fourier inversion and Cauchy–Schwarz give

$$\begin{aligned} |D_w(x) - \alpha^3| &\leq \sum_{\chi \neq 1} |\widehat{w}(\chi)|^3 \\ &\leq \left(\sum_{\chi \neq 1} |\widehat{w}(\chi)|^4 \right)^{1/2} \left(\sum_{\chi \neq 1} |\widehat{w}(\chi)|^2 \right)^{1/2} \\ &= u^2 \sqrt{\mu - \alpha^2}. \end{aligned}$$

Now use $\mu - \alpha^2 \leq \alpha(1 - \alpha)$ and the definition of λ . □

The exponent $\alpha^{5/4}$ in the main theorem is already visible here: controlling the triple link relative to its mean requires $u^2 \sqrt{\alpha}$ to be small compared with α^3 . This explains the sufficient exponent used by this proof; it does not establish its necessity.

II.255 Weighted modal differences without discarding fibers

Source 56, delivered Section 4.

For each h with $r(h) > 0$, let ν_h be the probability distribution of

$$\phi(x + h) - \phi(x)$$

when x is sampled with weights $w(x)w(x + h)$. Choose a mode $q(h) \in H$. When $r(h) = 0$, set $q(h) = 0$. Put

$$b(h) = 1 - \nu_h(q(h)), \quad c(h) = 1 - \sum_{v \in H} \nu_h(v)^2,$$

with both values zero on zero-mass fibers. By Lemma II.253.2, $0 \leq b(h) \leq c(h) \leq 1$. The total bad-pair mass is

$$B = \mathbb{E}_h r(h) b(h) = \mathbb{E}_{x,h} w(x) w(x + h) \mathbf{1}_{\phi(x+h) - \phi(x) \neq q(h)}. \quad (\text{II.255.1})$$

Lemma II.255.1 (Collision identity and weighted interpolation). *Writing $D = \mathcal{D}_w(\phi) = \varepsilon(\alpha^4 + u^4)$, one has*

$$D = \mathbb{E}_h r(h)^2 c(h) \quad (\text{II.255.2})$$

and

$$\alpha^2 B \leq D + u^2 \sqrt{D}. \quad (\text{II.255.3})$$

In particular,

$$\frac{B}{\alpha^2} \leq \beta := \varepsilon(1 + \kappa) + s \sqrt{\varepsilon(1 + \kappa)}. \quad (\text{II.255.4})$$

Proof. A quadruple with the same difference on two ordered pairs can be written as $(x + h, y, x, y + h)$. It is respected precisely when the two differences of ϕ coincide. The weight of such pairs at fixed h is $r(h)^2$, which proves (II.255.2). Since $b \leq c$,

$$\mathbb{E}_h r(h)^2 b(h) \leq D.$$

Insert $\alpha^2 = r(h) - (r(h) - \alpha^2)$ into $\alpha^2 B$:

$$\begin{aligned}\alpha^2 B &= \mathbb{E}_h r(h)^2 b(h) - \mathbb{E}_h (r(h) - \alpha^2) r(h) b(h) \\ &\leq D + \left(\mathbb{E}_h (r(h) - \alpha^2)^2 \right)^{1/2} \left(\mathbb{E}_h r(h)^2 b(h)^2 \right)^{1/2} \\ &\leq D + u^2 \sqrt{D}.\end{aligned}$$

The last step uses $b^2 \leq b$ and Lemma II.254.1. Dividing by α^4 gives (II.255.4). $\square$

The absence of a threshold on $r(h)$ is important. Selecting only shifts with $r(h)$ above a cutoff would introduce a further exceptional-set loss. Here both the original test and the selected difference map retain their actual pair multiplicities.

Lemma II.255.2 (Triangle transfer to an ambient additive test). *Let*

$$R_q = \left\{ (h, k) \in G^2 : q(h+k) \neq q(h) + q(k) \right\}, \quad \rho = \frac{|R_q|}{N^2}, \quad v = \frac{\sqrt{V}}{\alpha^3}.$$

Then

$$\rho \leq 3\beta + v\sqrt{\rho} \tag{II.255.5}$$

and therefore

$$\rho \leq \left(\frac{v + \sqrt{v^2 + 12\beta}}{2} \right)^2 \leq 6\beta + v^2 \leq 6\beta + 3\lambda^2. \tag{II.255.6}$$

Proof. For any $(h, k) \in R_q$, a triangle $(x, x+h, x+h+k)$ cannot have all three differences equal to $q(h), q(k), q(h+k)$. Thus at least one of its three edges is a bad pair in (II.255.1). An erroneous ordered pair extends to triangles with third-vertex total weight exactly α . This holds for each of the three edge positions, by bijective changes of variables. Hence

$$\mathbb{E}_{h,k} T(h, k) \mathbf{1}_{R_q}(h, k) \leq 3\alpha B. \tag{II.255.7}$$

On the other hand, Cauchy–Schwarz and (II.254.5) give

$$\alpha^3 \rho \leq \mathbb{E}_{h,k} T(h, k) \mathbf{1}_{R_q}(h, k) + \sqrt{V} \sqrt{\rho}.$$

Divide by α^3 and apply (II.255.4). Solving the quadratic inequality in $\sqrt{\rho}$ gives the first bound in (II.255.6). Alternatively, $v\sqrt{\rho} \leq (\rho + v^2)/2$ gives $\rho \leq 6\beta + v^2$. Finally $1 + 2\alpha \leq 3$ in (II.254.7). $\square$

The square-root expression is retained for data-dependent applications. The coarser last expression is convenient for a short rational verification of the universal constants.

II.256 Self-correction and coarse affine reconstruction

Source 56, delivered Section 5.

We include the classical self-correction argument in full because it fixes the threshold, works for arbitrary abelian targets, and makes the subsequent constant calculation independent of a black-box stability theorem. The method is the homomorphism self-correction mechanism of Blum–Luby–Rubinfeld [49]; it is not claimed as a new result.

Lemma II.256.1 (Ambient homomorphism self-correction). *Let $q : G \rightarrow H$ have additive-test error*

$$\rho = \mathbb{P}_{x,y}(q(x+y) \neq q(x) + q(y)) < \frac{1}{6}.$$

There is a homomorphism $L : G \rightarrow H$ such that

$$\tau := \mathbb{P}_x(q(x) \neq L(x)) \leq \frac{\rho}{1 - 2\rho}. \quad (\text{II.256.1})$$

For every x , the value $L(x)$ is the unique mode of $q(x+y) - q(y)$ as y ranges uniformly over G ; that mode has mass at least $1 - 2\rho$.

Proof. Fix x and write $d_x(y) = q(x+y) - q(y)$. For independent uniform y, z , the two identities

$$q(x+y+z) = q(x+y) + q(z), \quad q(y+z) = q(y) + q(z)$$

fail with combined probability at most 2ρ . When both hold, $d_x(y+z) = d_x(y)$. Since $(y, y+z)$ is a uniformly distributed independent pair, the collision probability of the distribution of d_x is at least $1 - 2\rho$. Lemma II.253.2 supplies a mode of at least this mass. It is unique because $1 - 2\rho > 1/2$. Define that mode to be $L(x)$.

The exact identity

$$d_{x+z}(y) = d_x(y) + d_z(x+y)$$

holds for all x, z, y . For fixed x, z , each of the three quantities equals its corresponding mode with probability at least $1 - 2\rho$ over y . The union bound leaves a positive probability of simultaneous agreement because $6\rho < 1$. It follows that $L(x+z) = L(x) + L(z)$ for every x, z .

If $q(x) \neq L(x)$, the conditional probability that $d_x(y) = q(x)$ is at most 2ρ . Thus the additive test fails, conditional on this x , with probability at least $1 - 2\rho$. Averaging over such x yields $\rho \geq (1 - 2\rho)\tau$, proving (II.256.1). $\square$

Lemma II.256.2 (Return to the weighted domain). *Suppose q has been selected as in Section II.255, and a homomorphism L differs from q on an ambient proportion τ . Then there is $c \in H$ such that*

$$t_w(\phi, L + c) \leq \beta + \tau + s\sqrt{\tau}. \quad (\text{II.256.2})$$

Proof. Let $S = \{h : q(h) \neq L(h)\}$. A weighted pair fails the relation $\phi(x+h) - \phi(x) = L(h)$ only if it is a bad pair for q or its shift lies in S . The total bad-pair mass is at most

$$B + \mathbb{E}_h r(h) \mathbf{1}_S(h) \leq B + \alpha^2 \tau + u^2 \sqrt{\tau}.$$

Here the last step is Cauchy–Schwarz applied to $(r - \alpha^2) \mathbf{1}_S$. After division by α^2 , the two endpoints of a weighted pair are independently distributed with mass function $w(x)/(N\alpha)$. Therefore the collision defect of $\psi(x) = \phi(x) - L(x)$ under that probability measure is at most $\beta + \tau + s\sqrt{\tau}$. Choose a mode c of ψ and apply Lemma II.253.2 once more. $\square$

Proposition II.256.3 (Parameterized coarse repair). *Compute β from (II.255.4) and $v = \sqrt{V}/\alpha^3$. Set*

$$\rho_* = \left(\frac{v + \sqrt{v^2 + 12\beta}}{2} \right)^2.$$

If $\rho_* < 1/6$, set $\tau_* = \rho_*/(1 - 2\rho_*)$. There is an affine map Φ with

$$t_w(\phi, \Phi) \leq \beta + \tau_* + s\sqrt{\tau_*}. \quad (\text{II.256.3})$$

One may replace ρ_* throughout by any larger upper bound below $1/6$, in particular by $6\beta + 3\lambda^2$ when that quantity is below $1/6$.

Proof. Combine Lemmas II.255.2, II.256.1, and II.256.2. The functions $x/(1 - 2x)$ and $x + s\sqrt{x}$ are increasing on their relevant nonnegative domains. $\square$

This proposition is already useful outside the fixed numerical range of Theorem F.61.1. It separates two tasks: establishing a coarse approximation from triangle sampling, and improving that approximation using the local error profile. Merely estimating the expression in (II.256.3) as $\varepsilon \rightarrow 0$ would leave a nonzero uniformity-dependent error. The next section is what removes that error.

II.257 The one-error profile and linear amplification

Source 56, delivered Section 6.

Fix an affine map $\Phi = L + c$. Subtracting it from ϕ does not alter which additive quadruples are respected. We may therefore suppose that $\Phi = 0$. Let

$$E = \{x \in \text{supp } w : \phi(x) \neq 0\}, \quad v_E(x) = w(x)\mathbf{1}_E(x), \quad b = \mathbb{E}v_E = \alpha t.$$

The letter v_E denotes a marked weight, not the normalized triangle standard deviation from Section II.256.

Lemma II.257.1 (Two marked vertices). *Place v_E at any two distinct coordinate positions in the ordered equation $a + b = c + d$, and place w at the other two. The resulting ambient-normalized weighted count M satisfies*

$$M \leq b^2\alpha^2 + u^2b^{3/2}. \quad (\text{II.257.1})$$

Proof. There are two types of positions. If the marked variables lie on the same side of the equation, the Fourier expression is

$$\sum_{\chi} \widehat{v_E}(\chi)^2 \overline{\widehat{w}(\chi)^2}.$$

If they lie on opposite sides, it is

$$\sum_{\chi} |\widehat{v_E}(\chi)|^2 |\widehat{w}(\chi)|^2.$$

In either case the trivial character contributes $b^2\alpha^2$ and the absolute value of the remaining contribution is at most

$$\begin{aligned} \sum_{\chi \neq 1} |\widehat{v_E}(\chi)|^2 |\widehat{w}(\chi)|^2 &\leq \left(\sum_{\chi \neq 1} |\widehat{v_E}(\chi)|^4 \right)^{1/2} u^2 \\ &\leq u^2 b^{3/2}, \end{aligned}$$

by (II.253.2). No division by two in G is used. The argument is therefore valid in groups with two-torsion as well. $\square$

Proposition II.257.2 (Local defect profile). *For every affine comparison map and its relative error t , the inequalities (F.61.8) and (F.61.9) hold.*

Proof. For an ordered additive quadruple let $K_i \in \{0, 1\}$ indicate that its i th coordinate lies in E . A quadruple with exactly one marked coordinate is necessarily violated: the three zero residual labels cannot cancel the single nonzero one. Pointwise on four binary indicators,

$$\mathbf{1}_{K_1+K_2+K_3+K_4=1} \geq \sum_{i=1}^4 K_i - 2 \sum_{1 \leq i < j \leq 4} K_i K_j. \quad (\text{II.257.2})$$

Indeed, for $r = 0, 1, 2, 3, 4$ marked incidences, the right side is $r - 2\binom{r}{2}$, respectively $0, 1, 0, -3, -8$.

For each fixed coordinate position, the weighted count with that coordinate marked is

$$\mathbb{E}_x v_E(x) D_w(x) \geq b\alpha^3(1 - \lambda\sqrt{1 - \alpha})$$

by Lemma II.254.3. The six two-marked counts are bounded by Lemma II.257.1. Multiplying (II.257.2) by the quadruple weight and averaging therefore gives

$$\mathcal{D}_w(\phi) \geq 4b\alpha^3(1 - \lambda\sqrt{1 - \alpha}) - 12b^2\alpha^2 - 12u^2b^{3/2}.$$

Substitute $b = \alpha t$, $u^2 = \lambda\alpha^{5/2}$, and $\mathcal{D}_w(\phi) = \varepsilon\alpha^4(1 + \lambda^2\alpha)$ to obtain (F.61.8).

Conversely, a violated quadruple must have at least one marked coordinate. The union bound and the upper estimate in Lemma II.254.3 yield

$$\mathcal{D}_w(\phi) \leq 4b\alpha^3(1 + \lambda\sqrt{1 - \alpha}),$$

which is (F.61.9). □

Repeated vertices cause no difficulty. The K_i count coordinate incidences, not distinct elements of E . If the same erroneous point occurs twice it contributes two incidences, and the exact-one assertion is not applied to it. This observation prevents an implicit nondegeneracy assumption.

II.257.1 A sharper form retaining the marked energy

The proof actually bounds each two-marked term by

$$b^2\alpha^2 + u^2\sqrt{\mathcal{Q}(v_E) - b^4}.$$

Thus, when the error support itself has small centered energy, the $t^{3/2}$ correction in (F.61.8) can be reduced. One fully explicit density-only variant is

$$\varepsilon(1 + \lambda^2\alpha) \geq 4(1 - \lambda\sqrt{1 - \alpha})t - 12t^2 - 12\lambda t^{3/2}\sqrt{1 - \alpha t}, \quad (\text{II.257.3})$$

because (II.253.2) gives $\mathcal{Q}(v_E) - b^4 \leq b^3(1 - b)$. We do not need this small additional improvement for the main theorem.

II.257.2 Why a coarse theorem suffices

The right side of (F.61.8) is linear and positive for sufficiently small t , with a coefficient close to four when λ is small. Thus a domain-dependent coarse estimate placing t in that region immediately becomes a bound proportional to ε . This is the central bootstrapping step. It avoids demanding that every intermediate approximation already have error $O(\varepsilon)$.

II.258 Proof of the main theorem and optimal first-order coefficient

Source 56, delivered Section 7.

II.258.1 A rational constant calculation

Assume (F.61.5). Then

$$\lambda \leq \frac{1}{32}, \quad \kappa \leq \frac{1}{1024}, \quad s \leq \frac{1}{32}, \quad \sqrt{\varepsilon} \leq \frac{1}{32}, \quad \sqrt{1+\kappa} \leq \frac{33}{32}.$$

Consequently (II.255.4) implies

$$\beta \leq \frac{1025}{1048576} + \frac{1056}{1048576} = \frac{2081}{1048576} < \frac{1}{480}. \quad (\text{II.258.1})$$

By Lemma II.255.2,

$$\rho \leq 6\beta + 3\lambda^2 < \frac{1}{80} + \frac{3}{1024} = \frac{79}{5120} < \frac{1}{64} < \frac{1}{6}. \quad (\text{II.258.2})$$

Apply Lemma II.256.1. Since

$$\tau \leq \frac{\rho}{1-2\rho} < \frac{1}{62} < \frac{1}{60}, \quad \sqrt{\tau} < \frac{1}{7},$$

Lemma II.256.2 gives an affine map with

$$t < \frac{1}{480} + \frac{1}{60} + \frac{1}{224} = \frac{13}{560} < \frac{1}{32}. \quad (\text{II.258.3})$$

We emphasize that no limit argument is being used; all constants are finite and the inequalities apply also when $\varepsilon = 0$.

For $t \leq 1/32$, the lower profile implies

$$\begin{aligned} \varepsilon(1+\kappa) &\geq t \left[4(1-\lambda\sqrt{1-\alpha}) - 12t - 12\lambda\sqrt{t} \right] \\ &\geq t \left[4\left(1 - \frac{1}{32}\right) - \frac{12}{32} - \frac{12}{32 \cdot 4} \right] = \frac{109}{32}t. \end{aligned}$$

Here $\sqrt{t} \leq 1/4$ is a convenient weaker bound than $1/\sqrt{32}$. Since

$$\frac{109}{32} > 3 \left(1 + \frac{1}{1024} \right),$$

we have established

$$t \leq \frac{\varepsilon}{3}. \quad (\text{II.258.4})$$

In particular, $\varepsilon = 0$ forces $t = 0$.

For $t > 0$, divide the lower profile by t and then use $12t \leq 4\varepsilon$ and $12\lambda\sqrt{t} \leq 4\sqrt{3}\lambda\sqrt{\varepsilon}$. This proves (F.61.10); the same conclusion holds trivially when $t = 0$. Its denominator is at least

$$4 \left(1 - \frac{1}{32} \right) - \frac{4}{1024} - 4 \cdot \frac{7}{4} \cdot \frac{1}{32} \cdot \frac{1}{32} = \frac{3957}{1024},$$

where $\sqrt{3} < 7/4$. Therefore

$$t \leq \frac{1025}{3957}\varepsilon \leq \frac{13}{50}\varepsilon. \quad (\text{II.258.5})$$

The constant inequality $1025/3957 < 13/50$ is strict; the displayed error bound is nonstrict so as to include zero error.

II.258.2 Separation and uniqueness

Lemma II.258.1 (Weighted affine separation). *Suppose $\sigma \leq \alpha$. Distinct affine maps $\Phi_1, \Phi_2 : G \rightarrow H$ satisfy*

$$\frac{\mathbb{E}_x w(x) \mathbf{1}_{\Phi_1(x) \neq \Phi_2(x)}}{\alpha} \geq \frac{1 - \sigma/\alpha}{2}. \quad (\text{II.258.6})$$

Proof. Their equality set is either empty or a coset of the kernel of the difference of their linear parts. If the linear parts are equal but their constants are different, the equality set is empty. Otherwise the kernel is a proper subgroup K of index $m \geq 2$. Its annihilator $K^\perp$ has m characters, and orthogonality gives

$$\mathbf{1}_{x_0+K}(x) = \frac{1}{m} \sum_{\chi \in K^\perp} \chi(x - x_0).$$

It follows that the relative weighted mass of this coset is at most

$$\frac{1}{m} + \left(1 - \frac{1}{m}\right) \frac{\sigma}{\alpha} \leq \frac{1 + \sigma/\alpha}{2}.$$

Taking the complement proves the lemma. $\square$

Under the main hypothesis,

$$\frac{\sigma}{\alpha} \leq \frac{u}{\alpha} = \sqrt{\lambda} \alpha^{1/4} \leq \frac{1}{\sqrt{32}} < \frac{1}{5}.$$

Thus distinct affine maps have relative weighted distance greater than $2/5$. Two maps each satisfying (F.61.6) would have mutual distance at most $26\varepsilon/50 \leq 26/(50 \cdot 1024) < 2/5$, a contradiction. Together with (II.258.5), this completes the proof of Theorem F.61.1. Taking $w = \mathbf{1}_A$ gives Corollary F.61.2, and Proposition II.257.2 plus the bootstrap proves Theorem F.61.3.

II.258.3 The coefficient $1/4$ is asymptotically exact

Corollary II.258.2. *For any sequence of inputs satisfying (F.61.5) with actual errors $\varepsilon > 0$, if $\varepsilon \rightarrow 0$ and $\lambda \rightarrow 0$, then the repaired distances satisfy*

$$\frac{t}{\varepsilon} \rightarrow \frac{1}{4}.$$

Proof. The upper profile gives

$$\frac{t}{\varepsilon} \geq \frac{1 + \lambda^2 \alpha}{4(1 + \lambda \sqrt{1 - \alpha})}.$$

The refined upper bound (F.61.10) gives the other side. Both comparison expressions tend to $1/4$ uniformly in $0 < \alpha \leq 1$. $\square$

Example II.258.3 (A single erroneous label). Let $G = \mathbb{Z}/N\mathbb{Z}$, let $H = \mathbb{Z}/2\mathbb{Z}$, set $w = 1$, and change the zero map at one point to the nonzero element of H . Its distance from the zero affine map is $t = 1/N$. The two profiles with $\lambda = 0$ give

$$4t - 12t^2 \leq \varepsilon \leq 4t.$$

For N sufficiently large the main theorem applies, and affine separation shows that the zero map is the unique nearest affine map. Hence $t/\varepsilon \rightarrow 1/4$. For completeness the exact error is

$$\varepsilon = \frac{(N-1)(4N-8)}{N^3} = 4t - 12t^2 + 8t^3.$$

At every nonzero shift the label-difference distribution has multiplicities $N-2$ and 2 , which proves the formula. This elementary full-domain example is only a sharpness witness; sharper full-domain stability profiles are already present in the inspected repository.

The equality of the limiting coefficient concerns the joint limit $\lambda, \varepsilon \rightarrow 0$. At fixed positive λ , the present upper and lower bounds need not have the same first-order coefficient. Neither an exact extremal profile nor an extremizer classification at fixed uniformity is asserted.

II.259 Exact extension, spectral certificates, and testing

Source 56, delivered Section 8.

Corollary II.259.1 (Exact Freiman extension). *Under $u^4 \leq \alpha^5/1024$, any map $\phi : \text{supp } w \rightarrow H$ respecting every additive quadruple of positive weight extends to a unique affine map on G .*

Proof. Apply Theorem F.61.1 with $\varepsilon = 0$. Relative weighted distance zero means equality at every point with positive weight because the domain is finite. Uniqueness follows from Lemma II.258.1. $\square$

This is an explicit exact-extension consequence of the same repair theorem; no passage through arbitrarily small positive error parameters is required.

Corollary II.259.2 (A maximum-Fourier-coefficient certificate). *If*

$$\max_{\chi \neq 1} |\widehat{w}(\chi)| \leq \frac{\alpha^2}{32}, \tag{II.259.1}$$

then the domain hypothesis of Theorem F.61.1 holds.

Proof. Parseval yields

$$u^4 = \sum_{\chi \neq 1} |\widehat{w}(\chi)|^4 \leq \sigma^2(\mu - \alpha^2) \leq \sigma^2 \alpha \leq \frac{\alpha^5}{1024}.$$

$\square$

The certificate is sufficient, not equivalent to the U^2 condition. Direct use of the fourth moment can be considerably better when a few Fourier coefficients dominate the spectrum.

II.259.1 Working relative to a subgroup or a coset

Let $K \leq G$ and let the domain lie in $x_0 + K$. Translate by $-x_0$ and use probability averages on K , not on G . The theorem gives a map

$$x_0 + k \mapsto L_K(k) + c, \quad L_K \in \text{Hom}(K, H),$$

with the same estimates, provided the relative $U^2(K)$ hypothesis holds. It does not in general extend that affine map from $x_0 + K$ to all of G .

Example II.259.3 (An ambient-extension obstruction). Take $G = \mathbb{Z}/4\mathbb{Z}$, $K = \{0, 2\}$, and $H = \mathbb{Z}/2\mathbb{Z}$. The map on K defined by $\phi(0) = 0$, $\phi(2) = 1$ is a homomorphism on K and hence respects every additive quadruple there. It has no affine extension to G : its constant would have to be zero, while every homomorphism $L : G \rightarrow H$ satisfies $L(2) = 2L(1) = 0$. Thus Fourier uniformity in the correct ambient group is a substantive hypothesis, and relative extension must not be silently identified with extension to a larger group.

II.259.2 A four-query affine tester

Define

$$d_w(\phi, \text{Aff}(G, H)) = \min_{\Phi \in \text{Aff}(G, H)} t_w(\phi, \Phi).$$

The minimum exists: each candidate gives the mass of one of finitely many error subsets of G , and at least the constant maps exist.

Corollary II.259.4 (Soundness with a weighted quadruple sampler). *Assume $u^4 \leq \alpha^5/1024$. The test that samples a weighted additive quadruple and checks its label relation has perfect completeness and rejection probability*

$$\varepsilon \geq \min \left\{ \frac{1}{1024}, \frac{50}{13} d_w(\phi, \text{Aff}(G, H)) \right\}. \quad (\text{II.259.2})$$

For distinct affine maps the relative weighted code distance is greater than $2/5$ under the same hypothesis.

Proof. Affine maps respect every quadruple. If $\varepsilon \leq 1/1024$, Theorem F.61.1 gives $d_w \leq 13\varepsilon/50$. If $\varepsilon > 1/1024$, the other branch of (II.259.2) is immediate. The code-distance assertion was proved in Section II.258. $\square$

This is a statement about *four label queries per accepted configuration*. It does not promise four membership queries to produce that configuration. For instance, rejection sampling starts with independent uniform a, b, c , sets $d = a + b - c$, and accepts with probability $w(a)w(b)w(c)w(d)$. The expected number of ambient proposals is $1/Q(w)$, which may be of order α^{-4} . Given an efficient weighted quadruple sampler, repeating the test m times gives false acceptance probability at most $(1 - \varepsilon)^m \leq e^{-m\varepsilon}$.

II.259.3 A local-to-global interpretation

The conclusion can be viewed as a quantitative presentation theorem for G : labels on a sufficiently uniform weighted subset that almost respect the four-term presentation relations nearly arise from a global affine map. The code-distance estimate gives uniqueness, whereas triangle transfer gives existence. These are logically separate: separation by itself would not produce a nearby global map.

II.260 Explicit affine repair on quadratic level sets

Source 56, delivered Section 9.

This section turns the Fourier hypothesis into a concrete algebraic rank condition. The labels are tested for *affinity*, even though their domain is a quadratic level set. Repair of approximately quadratic labels is a different problem and is not proved here.

Fix an odd prime p , and write $e_p(t) = \exp(2\pi it/p)$. A quadratic polynomial on $\mathbb{F}_p^n$ has the form

$$Q(x) = x^T M x + \ell(x) + c,$$

with M symmetric. In this section its rank means the matrix rank of M . The linear and constant terms do not affect this rank.

Lemma II.260.1 (Quadratic Fourier bound). *For every quadratic Q with matrix rank R_0 and every linear functional ξ , one has*

$$|\mathbb{E}_x e_p(Q(x) - \xi(x))| \leq p^{-R_0/2}.$$

Proof. Absorb $-\xi$ into the linear part of Q . Squaring the absolute value of the average and introducing a difference h gives

$$|\mathbb{E}_x e_p(Q(x))|^2 = \mathbb{E}_h \mathbb{E}_x e_p(Q(x+h) - Q(x)).$$

The dependence on x in the phase is the linear form $2x^T M h$. Its average vanishes unless $Mh = 0$, because p is odd. The remaining terms have modulus at most one. Since the kernel of M has relative size p^{-R_0} , the squared absolute value is at most p^{-R_0} . $\square$

Theorem II.260.2 (Rank certificate independent of the test error). *Let $r \geq 1$ and let $Q_1, \dots, Q_r : \mathbb{F}_p^n \rightarrow \mathbb{F}_p$ be quadratic polynomials such that every nonzero linear combination has matrix rank at least R . For any $b_1, \dots, b_r \in \mathbb{F}_p$, put*

$$A = \{x : Q_1(x) = b_1, \dots, Q_r(x) = b_r\}.$$

Suppose

$$p^{R-4r} \geq 2^{14}. \tag{II.260.1}$$

Then A is nonempty. For every abelian group H and every map $\phi : A \rightarrow H$ with additive-quadruple error at most $0 \leq \varepsilon \leq 1/1024$, there is a unique affine map $\Phi : \mathbb{F}_p^n \rightarrow H$ agreeing with ϕ except on at most $13\varepsilon|A|/50$ points.

Proof. Orthogonality gives

$$\mathbf{1}_A(x) = p^{-r} \sum_{\theta \in \mathbb{F}_p^r} e_p \left(\sum_{j=1}^r \theta_j (Q_j(x) - b_j) \right). \tag{II.260.2}$$

The term $\theta = 0$ has mean p^{-r} . By Lemma II.260.1,

$$|\alpha - p^{-r}| \leq (1 - p^{-r})p^{-R/2} \leq p^{-R/2}. \tag{II.260.3}$$

At a nontrivial additive character, the $\theta = 0$ term in the Fourier transform vanishes. The same argument gives

$$\sigma \leq p^{-R/2}. \tag{II.260.4}$$

The rank hypothesis implies

$$p^{-R/2} \leq \frac{p^{-2r}}{128} \leq \frac{p^{-r}}{2},$$

so $\alpha \geq p^{-r}/2 > 0$. In particular,

$$\sigma \leq \frac{p^{-2r}}{128} \leq \frac{\alpha^2}{32}.$$

Corollary II.259.2 and then Theorem F.61.1 apply. Equivalently, a direct calculation is

$$\frac{\|\mathbf{1}_A - \alpha\|_{U^2}^4}{\alpha^5} \leq \frac{p^{-R}\alpha}{\alpha^5} \leq 16p^{4r-R} \leq \frac{1}{1024}.$$

□

An equivalent integer rank bound is

$$R \geq 4r + \lceil 14 \log_p 2 \rceil.$$

For example, the sufficient additional ranks are nine for $p = 3$, seven for $p = 5$, and five for $p = 7$, since 3^9 , 5^7 , and 7^5 exceed 2^{14} whereas the preceding powers do not. These constants are certificates from the displayed proof, not optimal-rank assertions. If $r = 0$, the domain is the whole vector space and the theorem applies without any rank requirement.

II.260.1 Nonvacuity and sharpness at small densities

Proposition II.260.3. *For every fixed odd prime p and positive integer r , there are domains with densities tending to p^{-r} , normalized uniformity $\lambda \rightarrow 0$, and maps with $\varepsilon \rightarrow 0$ whose nearest affine distance satisfies $t/\varepsilon \rightarrow 1/4$.*

Proof. For each M , use r disjoint blocks of M variables and put

$$Q_j(x) = \sum_{i=1}^M x_{j,i}^2, \quad G_M = \mathbb{F}_p^M, \quad A_M = \{Q_1 = \cdots = Q_r = 0\}.$$

Every nonzero combination has rank at least M . Thus

$$\alpha_M = p^{-r} + O(p^{-M/2}), \quad \lambda_M^2 \leq p^{-M} \alpha_M^{-4} \rightarrow 0.$$

In particular $|A_M| \rightarrow \infty$. Give every point label zero in $\mathbb{F}_p$ except one point, which receives label one. Its distance from the zero map is $t_M = 1/|A_M|$. The upper profile proves $\varepsilon_M \rightarrow 0$, and the lower profile proves $\varepsilon_M = 4t_M(1 + o(1)) > 0$ for large M . The main theorem and affine separation identify the zero map as the unique nearest affine map for large M : any other affine map disagrees with zero on more than $2/5$ of the domain and hence cannot improve its vanishing error. This proves the claim. □

[write] **7 October 2026; sources 54–59.** *Checked at the write* (exact rationals): the constant chain of Section II.258, $1025/1048576 + 1056/1048576 = 2081/1048576 < 1/480$, $6/480 + 3/1024 = 79/5120 < 1/64$, $(1/64)/(1 - 2/64) = 1/62 < 1/60$, $1/60 < 1/49$, $1/480 + 1/60 + 1/224 = 13/560 < 1/32$, $4(31/32) - 12/32 - 12/128 = 109/32 > 3 \cdot 1025/1024$, $4(31/32) - 4/1024 - 7/1024 = 3957/1024$ and $1025/3957 < 13/50$; and the ranks of Theorem II.260.2: $3^8 < 2^{14} \leq 3^9$, $5^6 < 2^{14} \leq 5^7$, $7^4 < 2^{14} \leq 7^5$, so the additional ranks are 9, 7 and 5 as stated.

Thus the limiting coefficient is not merely an artifact of the full-domain case. It remains optimal at arbitrarily small fixed limiting densities. The rank needed for the affine-repair corollary has no dependence on $\log(\varepsilon^{-1})$ or on ε^{-1} . This does not establish the analogous independence for quadratic-label repair.

II.261 Source 59's chapter: exact character separation at Section 7 [write]

Added in the tenth write, 7 October 2026.

The next five sections are source 59's Sections 2–6, on the extraction of a Freiman homomorphism of order k (Gowers's Lemma 7.5 and Corollary 7.6). *Logically they continue the chapters of sources 24 and 31 on Section 7* (Sections II.72 and II.114) and source 10's torsion budget; they are printed here because no number printed earlier may move.

What is proved. For a finite symmetric forbidden set with q sign orbits, the exact worst-case number of characters that separate it is $\lfloor \log_2(q+1) \rfloor$ at radius $1/5$, $\lfloor \log_3(2q+1) \rfloor$ at radius $1/7$ for odd-order defects, and $\lfloor \log_p((p-1)q+1) \rfloor$ at radius $1/(2p)$ for p -power orders, every rank sharp (Theorem II.263.4); the exact small-ball law of one character (Lemma II.262.2) and strict adaptive deletion (Lemma II.263.3) are the key. Weighted localization then gives exact Freiman restrictions of order k over arbitrary abelian targets, $|B|/(5k)^{\lfloor \log_2(q+1) \rfloor}$ in general (Theorem II.264.1), with finite-image rounding (Theorem II.264.3), and, from $|\Gamma - \Gamma| \leq C|\Gamma|$, $|B'| \geq C^{-(2k+1)\log_2(5k)}|B|$ independently of the target (Corollary II.265.2); this answers source 24's mixed-torsion question (item II.8). Every symmetric set is realized as the defect set of a graph at a prescribed order (Theorem II.265.4); the one-character problem contains the lonely runner conjecture, whose known bounds transfer (Proposition II.266.2). The separation architecture is Green and Tao's, the finite-field kernel Gowers and Milićević's, as source 59 credits; the universal exponent can exceed Gowers's $4k$, and source 10's prime-cyclic bound is not improved, as it says.

Relation to the formal statements (checked at HEAD fe7165a3c). `lemma_7_5` and `corollary_7_6` (Sections06_07.lean, lines 183 and 197) have the exact companions `lemma_7_5_holds` (Proofs07FreimanRestriction.lean, line 267) and `corollary_7_6_holds` (Proofs07AdditiveRestriction.lean, line 100), stated for prime cyclic groups; nothing in this chapter is formalized.

II.262 Exact character distributions and order-sensitive radii

Source 59, delivered Section 2.

Let H_0 be finitely generated. By the structure theorem, its character group is a product of a finite group and a torus; equip it with normalized Haar measure. We first record exactly what a random character does to a single element. This elementary detail controls every torsion issue that follows.

Lemma II.262.1 (Evaluation and extension). *If $f \in H_0$ has finite order n , evaluation $\chi \mapsto \chi(f)$ sends Haar measure to the uniform measure on $\mathbb{T}[n]$. If f has infinite order, it sends Haar measure to Haar measure on $\mathbb{T}$. Every character on a subgroup of an abelian group extends to the full group.*

Proof. First prove the extension assertion. Suppose χ is defined on $K \leq H$ and $x \notin K$. If there is a positive integer n with $nx \in K$, choose the least such n and choose $u \in \mathbb{T}$ with $nu = \chi(nx)$. The definition $\chi(k+mx) = \chi(k) + mu$ is well-defined on $K + \mathbb{Z}x$. If no such n exists, choose u arbitrarily and use the same definition. A maximal extension argument proves the assertion, using divisibility of $\mathbb{T}$.

Every character of $\langle f \rangle$ therefore extends to H_0 . Evaluation is consequently onto $\mathbb{T}[n]$ in the finite-order case and onto $\mathbb{T}$ in the infinite-order case. A continuous surjective homomorphism of

compact groups pushes normalized Haar measure to normalized Haar measure. Alternatively, the finite case follows from equal fiber sizes, and the infinite case follows directly by integrating over the torus coordinates of $\hat{H}_0$. $\square$

Lemma II.262.2 (Exact open-arc probability). *For $0 < \theta < 1/2$ and $n \in \{2, 3, \dots\} \cup \{\infty\}$, set*

$$\beta_n(\theta) = \begin{cases} (2\lceil n\theta \rceil - 1)/n, & n < \infty, \\ 2\theta, & n = \infty. \end{cases} \quad (\text{II.262.1})$$

If $\text{ord}(f) = n$, then

$$\mathbb{P}_\chi(\|\chi(f)\|_{\mathbb{T}} < \theta) = \beta_n(\theta).$$

Proof. For finite n , the residues at distance strictly less than θ are zero and the pairs $\pm j/n$ with $1 \leq j < n\theta$. There are $2\lceil n\theta \rceil - 1$ of them. For infinite order, the relevant open arc has length 2θ . $\square$

The strict arc convention is essential. For example,

$$\beta_5(1/4) = \frac{3}{5}, \quad \beta_5(1/5) = \frac{1}{5}, \quad \beta_6(1/5) = \frac{1}{2}.$$

Thus one cannot replace the exact distribution by continuous circle measure when the evaluated element has small finite order.

Proposition II.262.3 (Optimal uniform probability cutoffs). *For every integer $r \geq 2$,*

$$\beta_n\left(\frac{1}{3r-1}\right) \leq \frac{1}{r} \quad \text{for all finite } n \geq r \text{ and for } n = \infty. \quad (\text{II.262.2})$$

For every odd integer $r \geq 3$,

$$\beta_n\left(\frac{1}{3r-2}\right) \leq \frac{1}{r} \quad \text{for all odd finite } n \geq r \text{ and for } n = \infty. \quad (\text{II.262.3})$$

In each statement, the displayed radius is the largest radius for which that uniform probability bound holds over the stated class of finite orders.

For a prime p and $a \geq 1$,

$$\beta_{p^a}\left(\frac{1}{2p}\right) \leq \frac{1}{p}, \quad \beta_\infty\left(\frac{1}{2p}\right) = \frac{1}{p}. \quad (\text{II.262.4})$$

No larger radius gives this contraction uniformly over all p -power orders.

Proof. For (II.262.2), write $j = \lceil n/(3r-1) \rceil$. If $j = 1$, then $\beta_n = 1/n \leq 1/r$. If $j \geq 2$, then

$$n \geq (j-1)(3r-1) + 1, \quad n - r(2j-1) \geq (r-1)(j-2) \geq 0.$$

This proves the finite-order assertion, and $2/(3r-1) \leq 1/r$ handles infinite order. At order $n = 3r-1$, any increase of the radius admits the two nearest nonzero roots, making the probability at least $3/(3r-1) > 1/r$.

For (II.262.3), put $j = \lceil n/(3r-2) \rceil$. The case $j = 1$ is the same. For $j \geq 2$, the least odd integer strictly larger than $(j-1)(3r-2)$ is that number plus d , where $d = 1$ if j is odd and $d = 2$ if j is even. Hence

$$n - r(2j-1) \geq (r-2)(j-2) + d - 2 \geq 0.$$

For odd $j \geq 3$ the final expression is at least $r - 3$, and for even j it is nonnegative directly. Infinite order again satisfies the bound. The odd order $3r - 2$ proves radius optimality.

If p is odd, then $p^a/(2p)$ is a half-integer, so the finite probability in (II.262.4) equals $1/p$ exactly. For $p = 2$, it is $1/2$ at $a = 1$ and $1/2 - 1/2^a$ at $a \geq 2$. Finally $\beta_{p^a}(\theta) \rightarrow 2\theta$ as $a \rightarrow \infty$, which proves the last optimality assertion. $\square$

Taking $r = 2$ in (II.262.2) gives radius $1/5$ and contraction $1/2$. Taking $r = 3$ in (II.262.3) gives radius $1/7$ and contraction $1/3$. If the actual defects avoid all small orders, the larger values of r give additional options. Those options are sufficient estimates; no corresponding exact-rank claim is made here for every composite r .

II.263 Adaptive character separation and matching obstructions

Source 59, delivered Section 3.

Choose a representative set $D \subset F$ of the sign orbits and write $q = |D|$. All sets considered in this section are finite; the containing abelian group may be arbitrary. Characters may be chosen on a finitely generated subgroup and extended by Lemma II.262.1.

Theorem II.263.1 (Exact order-profile certificate). *Choose radii $0 < \theta_j < 1/2$, $1 \leq j \leq m$, and put*

$$P = \sum_{f \in D} \prod_{j=1}^m \beta_{\text{ord}(f)}(\theta_j). \quad (\text{II.263.1})$$

If $P < 1$, there are characters $\chi_1, \dots, \chi_m$ such that

$$\text{for every } f \in F, \quad \|\chi_j(f)\|_{\mathbb{T}} \geq \theta_j \text{ for some } j. \quad (\text{II.263.2})$$

For $q \geq 2$, the hypothesis $P \leq 1$ is sufficient. More generally, if $P < q$, some tuple of characters leaves at most $\lceil P \rceil - 1$ representatives satisfying all the strict reverse inequalities.

Proof. Take independent Haar-uniform characters and let X count representatives satisfying $\|\chi_j(f)\|_{\mathbb{T}} < \theta_j$ for every j . The exact probability formula gives $\mathbb{E}X = P$. All these inequalities hold on an open neighborhood of the zero character tuple, so $X = q$ on a set of positive Haar measure. If $P < q$ and $X \geq P$ everywhere, that neighborhood makes $\mathbb{E}X > P$, a contradiction. Thus $X < P$ somewhere; since X is integral, $X \leq \lceil P \rceil - 1$. The assertions follow. When $P < 1$ and $q = 1$, the elementary expectation argument also applies directly. $\square$

Remark II.263.2 (Why the endpoint needs an argument). An ordinary union bound with expectation one does not exclude a survivor at every character tuple. Here the bad events have a common open intersection of positive measure, on which at least two representatives survive. That extra overlap is what permits equality in the budget. For $q = 1$, a budget of one is not sufficient in general: a wide enough arc can contain the entire image of an odd-order element.

Lemma II.263.3 (Strict adaptive deletion). *Suppose every order occurring in F satisfies $\beta_n(\theta) \leq 1/b$, where $b \geq 2$ is an integer. Then F can be separated at radius θ by at most*

$$m_b(q) = \lfloor \log_b((b-1)q + 1) \rfloor \quad (\text{II.263.3})$$

characters, with $m_b(0) = 0$.

Proof. At a stage with $s > 0$ surviving representatives, let X count survivors after one new character. We have $\mathbb{E}X \leq s/b$. A positive-measure neighborhood of zero has $X = s > s/b$, so some character has $X < s/b$. Therefore

$$s_{\text{next}} \leq \left\lceil \frac{s}{b} \right\rceil - 1 = \left\lfloor \frac{s-1}{b} \right\rfloor. \quad (\text{II.263.4})$$

Writing $U(s) = (b-1)s + 1$, this implies $U(s_{\text{next}}) \leq U(s)/b$. Put $m = m_b(q)$. If the process has not terminated earlier, after m steps

$$U(s_m) \leq \frac{(b-1)q + 1}{b^m} < b.$$

Every positive integer s has $U(s) \geq b$, so $s_m = 0$. □

For perspective, m steps suffice exactly throughout the range

$$q \leq b + b^2 + \cdots + b^m.$$

Thus two universal characters handle as many as six sign orbits; two odd-order characters handle as many as twelve. The gain over taking a fixed number of independent characters is an adaptive integer gain, not a change in a continuous logarithmic approximation.

Theorem II.263.4 (Sharp universal, odd, and prime-primary ranks). *Let $F \subset H \setminus \{0\}$ be finite and symmetric, with q sign orbits.*

- (i) *With no restriction on H or F , radius $1/5$ is achievable with at most $m_2(q) = \lfloor \log_2(q+1) \rfloor$ characters.*
- (ii) *If every finite-order element of F has odd order, radius $1/7$ is achievable with at most $m_3(q) = \lfloor \log_3(2q+1) \rfloor$ characters.*
- (iii) *If every finite-order element of F has order a power of a fixed prime p , radius $1/(2p)$ is achievable with at most $m_p(q)$ characters.*

Each number of characters is best possible as a worst-case function of q , for every $q \geq 1$. In (i) and (ii), the radius cannot be increased while retaining that same rank bound for all q .

Proof. The upper bounds follow from Proposition II.262.3 and Lemma II.263.3.

For the rank lower bounds, fix p and $q \geq 1$, and put $m = m_p(q)$. Then

$$\frac{p^m - 1}{p - 1} \leq q < \frac{p^{m+1} - 1}{p - 1}.$$

Inside $V = \mathbb{F}_p^{m+1}$ fix an m -dimensional subspace U . From every one-dimensional subspace $L \leq U$, choose a nonzero representative v_L and include both v_L and $-v_L$ in F . Add such pairs from further distinct projective lines of V until there are exactly q sign orbits. The displayed inequalities ensure that this is possible, also for $p = 2$, when each pair is a singleton.

Characters of V identify with linear functionals into $\mathbb{F}_p$. Any $m-1$ characters have a nonzero common kernel on U . That kernel contains a projective line whose chosen representative belongs to F . Consequently $m-1$ characters fail at every positive radius. This proves the prime-primary lower bound for every q ; $p = 2$ and $p = 3$ prove the universal and odd-order lower bounds.

For universal radius sharpness, take $H = \mathbb{Z}/5\mathbb{Z}$ and $F = H \setminus \{0\}$. Here $q = 2$ and $m_2(q) = 1$. Every nontrivial character permutes $\mathbb{T}[5]$, so some forbidden element has distance $1/5$; the trivial character does worse. For the odd case, $H = \mathbb{Z}/7\mathbb{Z}$ and $F = H \setminus \{0\}$ have $q = 3$, $m_3(q) = 1$, and the analogous obstruction at $1/7$. $\square$

Remark II.263.5 (Three distinct optimality assertions). The theorem proves an exact rank for each q . It proves an optimal radius uniformly in q at those ranks in the first two cases. Proposition II.262.3 proves a pointwise probability threshold for prime powers. The last assertion does not imply that $1/(2p)$ is the largest attainable prime-primary separation radius. Nor does sharp separator rank by itself prove an optimal retained proportion in a Freiman restriction.

Example II.263.6 (Order profile versus a uniform class). Suppose the representative orders are 2 and 5. At radius $1/5$ the one-character budget is $1/2 + 1/5 = 7/10 < 1$, so one character suffices. If there are three representatives, one of order 3 and two of infinite order, then radius $1/7$ gives budget $1/3 + 2(2/7) = 19/21 < 1$. Exact profiles retain information that is discarded by replacing every survival probability by its largest allowed value. In the adaptive version, one may use the sharper bound $\lceil \sum_f \beta_{\text{ord}(f)}(\theta) \rceil - 1$ at each stage.

II.264 Weighted Freiman restrictions over arbitrary abelian targets

Source 59, delivered Section 4.

Let G, H be abelian groups, let $B \subset G$ be finite and nonempty, and let $\phi : B \rightarrow H$. Define Γ, F_k^0, F_k as in (F.64.1), where $k \geq 1$ is an integer. The group $H_0 = \langle \phi(B) \rangle$ is finitely generated, and $F_k \subset H_0$. Let $w : B \rightarrow [0, \infty)$ and write $w(B) = \sum_{x \in B} w(x)$.

Theorem II.264.1 (Weighted localization). *Suppose $\chi_1, \dots, \chi_m : H_0 \rightarrow \mathbb{T}$ separate F_k with radii $0 < \theta_j \leq 1/2$, in the sense of (II.263.2). Then some $B' \subset B$ satisfies*

$$w(B') \geq w(B) \prod_{j=1}^m \frac{\theta_j}{k}, \quad (\text{II.264.1})$$

and $\phi|_{B'}$ is a Freiman homomorphism of order k . Moreover, B can be partitioned into at most $\prod_j \lceil k/\theta_j \rceil$ subsets on each of which ϕ is such a homomorphism.

Proof. For each j , choose a uniformly translated half-open arc of length θ_j/k , independently in each circle coordinate. Retain x when every $\chi_j(\phi(x))$ lies in its corresponding arc. Every fixed x is retained with probability $\prod_j \theta_j/k$, so the expected retained weight is the right-hand side of (II.264.1). A translate with at least that weight exists.

Suppose $x_1 + \dots + x_k = y_1 + \dots + y_k$ with all $x_i, y_i \in B'$. Its label defect

$$h = \phi(x_1) + \dots + \phi(x_k) - \phi(y_1) - \dots - \phi(y_k)$$

belongs to F_k^0 . Lift each character value to the real half-open interval representing its chosen arc. Each of the two k -term sums lies in the same interval of length θ_j . Their difference therefore has absolute value strictly less than θ_j , so $\|\chi_j(h)\|_{\mathbb{T}} < \theta_j$ for every j . Separation forces $h = 0$.

For the partition assertion, divide circle j into $L_j = \lceil k/\theta_j \rceil$ equal half-open arcs. Each arc has length at most θ_j/k , and the same argument applies to every pulled-back product cell. Empty cells may be discarded. The half-open convention preserves strictness even when $1/L_j = \theta_j/k$. $\square$

The assertion also covers zero total weight: one can choose any nonempty cell, or a singleton, without affecting the inequality. If F_k is empty, there is no separation to perform and one takes $B' = B$.

Corollary II.264.2 (Explicit universal and torsion-sensitive retention). *Put $q = |F_k/\{\pm 1\}|$. An order- k Freiman restriction can retain at least the following fractions of any nonnegative weight:*

<i>Defect condition</i>	<i>Retained fraction</i>
<i>No restriction</i>	$(5k)^{-m_2(q)}$
<i>All finite defect orders odd</i>	$(7k)^{-m_3(q)}$
<i>All finite defect orders powers of p</i>	$(2pk)^{-m_p(q)}$

The exponent is zero when $q = 0$.

Proof. Apply Theorem II.263.4 and then Theorem II.264.1. □

II.264.1 Exact finite-image rounding

The continuous-arc loss is often wasteful on a finite character image. The following refinement keeps the finite arithmetic exactly.

Theorem II.264.3 (Finite-image gain). *In Theorem II.264.1, suppose $\chi_j(H_0)$ has finite order N_j . Its factor θ_j/k in (II.264.1) can be replaced by*

$$a_j = \frac{\lceil N_j \theta_j / k \rceil}{N_j}. \quad (\text{II.264.2})$$

If the image is infinite, keep the factor θ_j/k . All the improved factors can be multiplied, regardless of dependencies between the character coordinates.

Proof. Choose $t_j = \lceil N_j \theta_j / k \rceil$ consecutive points of the cyclic image and translate this block uniformly by that image. A fixed image point is included with probability t_j/N_j . The points in a block lift to a real interval of span

$$\frac{t_j - 1}{N_j} < \frac{\theta_j}{k}.$$

Thus the proof of exact Freiman behavior is unchanged. Make these translations independently in the ambient product of the coordinate images. They need not be induced by a common translation in H_0 ; every point still has the stated product inclusion probability. For a trivial image, $N_j = 1$ and its factor is one. □

Corollary II.264.4 (Finite exponent and finite-field fibers). *If H_0 has exponent E , the partition form of Corollary II.264.2 has at most $\min(E, dk)^m$ cells, where (d, m) is $(5, m_2(q))$, $(7, m_3(q))$, or $(2p, m_p(q))$ in the appropriate case.*

If H_0 has exponent p , an order- k restriction retains at least

$$p^{-m_p(q)} w(B). \quad (\text{II.264.3})$$

More strongly, let r be the number of distinct one-dimensional $\mathbb{F}_p$ -subspaces meeting F_k . Then q can be replaced by r in (II.264.3).

Proof. An exponent- E character has at most E values, so at most $\min(E, dk)$ equal arcs can be occupied in each coordinate.

For exponent p , a nontrivial character has image order p . At radius $1/(2p)$, formula (II.264.2) is exactly $1/p$ for every $k \geq 1$. Furthermore, a narrow-arc survivor is precisely a vector annihilated by the corresponding linear functional. Nonzero scalar multiples define the same annihilation event. Apply the adaptive argument to the r projective lines, whose individual survival probability is $1/p$, and take a common fiber of the selected functionals. $\square$

The linear-kernel method in this corollary is classical [237, Lemma 3.3]. Its role here is to show that the general circle formulation includes exact finite-field fibers, rather than unnecessarily charging a factor k for each character. The adaptive rank improves the rounding endpoint of the direct independent-kernel argument.

II.264.2 A projection formulation

Proposition II.264.5 (Exact local sections of an abelian quotient). *Let $\pi : E \rightarrow G$ be a homomorphism of abelian groups, let $A \subset E$ be finite, and put*

$$F = ((kA - kA) \cap \ker \pi) \setminus \{0\}.$$

If characters on $\langle A \rangle$ separate F at radii θ_j , the same retained-weight bound gives $A' \subset A$ on which π is a Freiman isomorphism of order k onto $\pi(A')$.

Proof. The localization proof shows that any equality of two k -term sums after applying π already holds in E . Forward preservation follows from the homomorphism property of π . If A' is nonempty, pad a one-term equality using any fixed point of A' to obtain injectivity; the empty case is vacuous. $\square$

This formulation allows non-split extensions: the ambient abelian group need not be presented as a product. The graph case is the special case of a projection whose restriction to the initial set is already injective.

II.265 Small doubling and comparison with Lemma 7.5

Source 59, delivered Section 5.

Lemma II.265.1 (Vertical injection; classical defect budget). *If $|\Gamma - \Gamma| \leq C|\Gamma|$, then*

$$Q_k := |F_k^0| \leq C^{2k+1}. \quad (\text{II.265.1})$$

Proof. The map

$$F_k^0 \times \Gamma \longrightarrow (k+1)\Gamma - k\Gamma, \quad (h, \gamma) \longmapsto (0, h) + \gamma$$

is injective. Indeed, the domain projection of its value determines the unique graph point γ , and then determines h . Therefore

$$Q_k |\Gamma| \leq |(k+1)\Gamma - k\Gamma|.$$

The classical Plünnecke–Ruzsa inequality says that $|A + B| \leq C|A|$ implies $|rB - sB| \leq C^{r+s}|A|$ for $r + s > 1$; see [36, Theorem 1.2]. Apply it with $A = \Gamma$, $B = -\Gamma$, $r = k$, and $s = k + 1$. Its conclusion concerns the full sum-difference set; no preliminary restriction of Γ is required. $\square$

The injection and this use of Plünnecke–Ruzsa are already present in the finite-field precedent and in sources 10 and 24 [237, 240]. We include them to make the quantitative interface explicit.

Corollary II.265.2 (Target-independent polynomial restriction). *Under the hypotheses of Lemma II.265.1, there is an order- k Freiman restriction satisfying*

$$|B'| \geq C^{-(2k+1)\log_2(5k)}|B|. \quad (\text{II.265.2})$$

If every finite-order nonzero defect has odd order, then

$$|B'| \geq C^{-(2k+1)\log_3(7k)}|B|. \quad (\text{II.265.3})$$

The same conclusions hold for any fixed nonnegative weight on B .

Proof. Universally, $q + 1 \leq Q_k$, so

$$m_2(q) \leq \log_2 Q_k \leq (2k + 1) \log_2 C.$$

In the odd-defect case, F_k has no involutions, so $2q + 1 = Q_k$ and $m_3(q) \leq \log_3 Q_k$. Substitute these estimates into Corollary II.264.2. For $C = 1$, the defect set is empty and the assertion still holds. $\square$

[write] **7 October 2026; sources 54–59.** *Source 24’s mixed-torsion question is answered* (its Question 7, item II.8): the bound holds for every abelian target, independently of its size, rank and invariant factors, with order-sensitive variants. Source 10’s prime-cyclic bound is not improved, as this source says. Source 61, placed later and not yet written in, records the same exponent $(2m + 1) \log_2(5m)$ found independently.

[write] **7 October 2026; sources 60–69.** *Checked at the eleventh write* (Section II.310): source 61 finds the exponent $(2m + 1) \log_2(5m)$ independently, in the partition form $|B|/(5mK^{(2m+1)\log_2(5m)})$ for some cell (its Theorem F.66.1); this corollary is sharper by the factor $5m$, and its defect-sensitive form by the separation rank $\lfloor \log_2(q + 1) \rfloor \leq \lceil \log_2 |D| \rceil$. Source 61’s pin precedes this source’s placement.

For mixed finite targets, the exact order of a defect is the least common multiple of its component orders. Formula (II.262.1) therefore uses its invariant-factor information without splitting the problem into primary components. The resulting estimate is independent of target cardinality and rank. It answers source 24’s mixed-torsion extension question in this explicit sense. A sharper bound for a particular group or a particular defect set can use Theorem II.263.1, finite-image rounding, or both.

For comparison, the exponents of C in several bounds are:

k	Gowers $4k$	Universal $(2k + 1) \log_2(5k)$	Odd $(2k + 1) \log_3(7k)$
2	8	16.610	12.011
4	16	38.897	27.298
8	32	90.473	62.289

The numerical entries are illustrative approximations. Gowers’s original statement has the additional prefactor $1/(8k)$ and a prime-cyclic target; the present rows have the broader target hypotheses stated above. Source 10’s prime-cyclic refinement is stronger still in that special setting. The actual Q_k and the finite-image corrections should be used before this coarse table whenever they are available.

II.265.1 An explicitly attributed Balog–Szemerédi–Gowers corollary

For a finite subset A of an abelian group, Reiher–Schoen prove that $E(A) \geq \gamma|A|^3$ and $0 < \eta < 1/2$ imply the existence of $A_1 \subset A$ with

$$|A_1| \geq (1 - \eta)\gamma^{1/2}|A|, \quad |A_1 - A_1| \leq 2^{33}\eta^{-9}\gamma^{-4}|A_1|. \quad (\text{II.265.4})$$

This is their Theorem 1.2 [12]; it is an external input, not a new estimate in this article.

Corollary II.265.3 (From graph energy to exact restrictions). *Suppose $E(\Gamma) \geq \gamma|B|^3$, with $0 < \gamma \leq 1$. Define*

$$A_k = (2k + 1) \log_2(5k).$$

Then, for every $0 < \eta < 1/2$, an order- k Freiman restriction satisfies

$$|B'| \geq (1 - \eta) 2^{-33A_k} \eta^{9A_k} \gamma^{1/2+4A_k} |B|. \quad (\text{II.265.5})$$

If the target has no nonzero finite-order elements of even order, one may instead take $A_k = (2k + 1) \log_3(7k)$. It is enough that all defects of the original graph have the odd-order property.

Proof. Apply (II.265.4) to Γ . Its subset is still a graph, now over some B_1 . Apply Corollary II.265.2 with $C = 2^{33}\eta^{-9}\gamma^{-4}$ and multiply the retained fractions. Defects of a restricted graph form a subset of the original defect set, so the odd-order condition is inherited. $\square$

For example, fixing $\eta = 1/4$ writes the universal factor as

$$\frac{3}{4} 2^{-51A_k} \gamma^{1/2+4A_k}.$$

This gives a fully explicit general-target analogue of the energy-to-Freiman step. It should not be described as a new record for Balog–Szemerédi–Gowers or as uniformly stronger than the project’s specialized prime-cyclic Corollary 7.6 refinements.

II.265.2 Every symmetric forbidden set is a graph defect set

The sharp separation examples are not excluded by requiring the forbidden set to arise from a graph. In fact, arbitrary symmetric forbidden sets admit the following exact realization. The construction also shows why sharp character rank does not establish sharp retained mass.

Theorem II.265.4 (Exact realization with prescribed Freiman order). *Let H be an abelian group, let $F \subset H \setminus \{0\}$ be finite, nonempty and symmetric, and let $q = |F/\{\pm 1\}|$. For each $k \geq 2$ there are an abelian group G , a set $B \subset G$ of size $2kq$, and a map $\phi : B \rightarrow H$ whose graph Γ satisfies*

$$\{h \neq 0 : (0, h) \in k\Gamma - k\Gamma\} = F.$$

One may take $G = \mathbb{Z}^{(2k-1)q}$, or the finite group $G = C_L^{(2k-1)q}$ for any integer $L > 2k$.

There is a partition $B = B_1 \sqcup \cdots \sqcup B_q$, with $|B_i| = 2k$, such that a subset $B' \subset B$ supports an exact Freiman restriction of order k if and only if it omits at least one point of every B_i . Consequently the largest such subset has size $(2k - 1)q$. For arbitrary nonnegative weights w on B , the largest retained weight is exactly

$$w(B) - \sum_{i=1}^q \min_{x \in B_i} w(x).$$

Moreover, ϕ itself is a Freiman homomorphism of every order $1 \leq j < k$.

Proof. Choose sign representatives $f_1, \dots, f_q$ of F and a basis $e_{i,j}$ of $\mathbb{Z}^{(2k-1)q}$, indexed by $1 \leq i \leq q$, $1 \leq j \leq 2k-1$. In the i th coordinate block put

$$v_{i,j} = e_{i,j} \quad (1 \leq j < 2k), \quad v_{i,2k} = \sum_{j=1}^k e_{i,j} - \sum_{j=k+1}^{2k-1} e_{i,j}.$$

These $2kq$ vectors are distinct. Set $B_i = \{v_{i,1}, \dots, v_{i,2k}\}$, put $\phi(v_{i,1}) = f_i$, and give all other vertices label zero. The relation

$$\sum_{j=1}^k v_{i,j} = \sum_{j=k+1}^{2k} v_{i,j}$$

has label defect f_i , so every element of F occurs.

Conversely, consider an equality of two sums of k vertices and let $a_{i,j}$ be the multiplicity on the left minus the multiplicity on the right. Then $\sum_{i,j} \max(a_{i,j}, 0) \leq k$. Coordinate comparison shows that the only integer relations in block i are

$$(a_{i,1}, \dots, a_{i,2k}) = t_i (\underbrace{1, \dots, 1}_k, \underbrace{-1, \dots, -1}_k) \quad (t_i \in \mathbb{Z}).$$

Hence $k \sum_i |t_i| \leq k$. Either every t_i is zero, in which case the two multisets are identical, or exactly one t_i equals 1 or -1 . In the latter case the relation uses every vertex of its block, once on the prescribed side, and has defect f_i or $-f_i$. This proves the asserted defect set and the necessary and sufficient condition on B' . Omitting a vertex of minimum weight from each block gives the weighted optimum. Replacing the two sums of length k by sums of length $j < k$ gives $k \sum_i |t_i| \leq j$, forcing every $t_i = 0$.

For the finite-domain version reduce the vectors modulo L . They remain distinct because their coordinate differences have absolute value at most $2 < L$. In a relation between two k -term sums, each coordinate of the corresponding integer sum is $a_{i,j} + a_{i,2k}$ or $a_{i,j} - a_{i,2k}$, and has absolute value at most $2k < L$. If that coordinate vanishes modulo L , it therefore vanishes as an integer. Thus the relations relevant to the proof are unchanged, and all conclusions persist. $\square$

For $F = \emptyset$, a one-point constant graph provides a realization. Theorem II.265.4 upgrades all the every- q separation-rank obstructions to graph-defect examples at any prescribed $k \geq 2$. Nevertheless their optimal retained proportion is $1 - 1/(2k)$, independently of F .

The construction also has an explicit additive cost. Writing $M = 2kq$,

$$|\Gamma - \Gamma| = M(M-1) + 1, \quad \frac{|\Gamma - \Gamma|}{|\Gamma|} = M - 1 + \frac{1}{M}.$$

Indeed, for $k > 2$ every two-term relation in the domain is trivial. For $k = 2$, the only nontrivial two-term relations have nonzero label defects and therefore are not relations in Γ . Thus every nonzero ordered difference of graph points has a unique representation; the zero difference accounts for the remaining one element. Bare graph realizability imposes no restriction on F , but requiring small difference doubling is a substantive additional condition.

II.266 One-character restrictions and lonely runners

Source 59, delivered Section 6.

The logarithmic rank is necessary in general because low-order torsion can force a common kernel. When the defects have infinite order, the relevant small-ball probabilities are continuous and a single character is enough for a useful linear bound.

Proposition II.266.1 (A one-character bound from the actual orders). *Let F have $q \geq 1$ sign representatives D and put*

$$T(F) = \sum_{\substack{f \in D \\ \text{ord}(f) < \infty}} \frac{1}{\text{ord}(f)}.$$

If $T(F) < 1$, then F can be separated by one character at radius

$$\theta = \frac{1 - T(F)}{2q}. \quad (\text{II.266.1})$$

Consequently, for a graph with this defect set, a weighted order- k Freiman restriction retains at least

$$\frac{1 - T(F)}{2kq} w(B). \quad (\text{II.266.2})$$

In particular, when every defect has infinite order, the fraction is $1/(2kq)$.

Proof. For finite n and $0 < \theta < 1/2$,

$$\beta_n(\theta) = \frac{2\lceil n\theta \rceil - 1}{n} < 2\theta + \frac{1}{n}.$$

If at least one finite-order defect occurs, the sum of the survival probabilities at (II.266.1) is strictly less than $2q\theta + T(F) = 1$. The profile theorem gives separation.

If all orders are infinite and $q \geq 2$, the radius $1/(2q)$ makes the sum exactly one, so the endpoint assertion of Theorem II.263.1 applies. If $q = 1$, define a character on the infinite cyclic subgroup generated by its representative by sending that representative to $1/2$, then extend it. This separate argument is needed because the open radius- $1/2$ arc has full Haar measure on the circle. In every case, apply Theorem II.264.1. $\square$

This is a sufficient torsion budget, not an exact criterion for one-character separation. It is useful when finite-order defects have large orders or are sparse. The exact profile in Theorem II.263.1 is normally stronger than the estimate $2\theta + 1/n$ used to obtain this compact formula.

II.266.1 An exact connection to the lonely runner constant

For $q \geq 1$, define

$$\Lambda(q) = \inf_{\substack{v_1, \dots, v_q \in \mathbb{Z}_{>0} \\ \text{distinct}}} \max_{t \in \mathbb{T}} \min_{1 \leq i \leq q} \|tv_i\|_{\mathbb{T}}. \quad (\text{II.266.3})$$

The integer lonely runner conjecture is the assertion $\Lambda(q) = 1/(q+1)$ for every q ; see [239]. We need no assumption of that conjecture.

Proposition II.266.2 (Transfer of the lonely runner constant). *If every element of F has infinite order, then one character separates F at radius $\Lambda(q)$. A weighted Freiman restriction therefore retains at least $\Lambda(q)w(B)/k$. Conversely, a uniform one-character assertion at radius $1/(q+1)$ for all such sets would imply the integer lonely runner conjecture.*

Proof. Work in a finitely generated subgroup H_0 containing F and the graph labels. Write $H_0/T_0 \cong \mathbb{Z}^d$, where T_0 is its finite torsion subgroup. Every defect has a nonzero image in this quotient. Choose an integer linear functional on $\mathbb{Z}^d$ which is nonzero on every such image and distinguishes the distinct images up to sign. This amounts to avoiding finitely many proper rational hyperplanes; such an integer functional exists, for example by evaluating at $(1, M, \dots, M^{d-1})$ for an integer M outside the finitely many roots of the resulting nonzero polynomials.

The absolute values of the projected defects give $r \leq q$ distinct positive integer speeds. Some $t \in \mathbb{T}$ realizes a minimum distance at least $\Lambda(r) \geq \Lambda(q)$. The maximum for a fixed finite set of speeds is attained by compactness; monotonicity follows by adjoining extra speeds. Composing the functional with multiplication by t gives the required character. Elements whose free images coincide or are opposite produce the same speed and cause no problem, even if the original elements differ by torsion.

For the converse take $H = \mathbb{Z}$ and $F = \{\pm v_1, \dots, \pm v_q\}$. Every character has the form $n \mapsto tn$, so the asserted separation is exactly the lower bound in the conjecture. The reverse inequality $\Lambda(q) \leq 1/(q+1)$ follows from speeds $1, 2, \dots, q$: among $0, t, \dots, qt$ two points have circular distance at most $1/(q+1)$, whereas $t = 1/(q+1)$ attains that value. $\square$

The elementary argument above gives $\Lambda(q) \geq 1/(2q)$. It is not the strongest known general lower bound. Bedert's Theorem 1.3 gives

$$\Lambda(q) \geq \frac{1}{2q} + q^{-5/3-o(1)} \quad (q \rightarrow \infty) \quad (\text{II.266.4})$$

in the cited preprint [238]. Proposition II.266.2 immediately transfers that additional term, divided by k , to infinite-order graph defects. The improvement in (II.266.4) is Bedert's, and the asymptotic $o(1)$ is not an explicit finite- q constant supplied by this article.

II.266.2 An explicit Riesz-product refinement under dissociativity

There is a complementary route for a structurally restricted forbidden set. The following statement extends the elementary Riesz-product calculation in Bedert's Lemma 4.1 [238] to arbitrary abelian groups and to a specified dissociativity order. Its method is classical; the explicit formula is useful as a reusable interface.

Definition II.266.3. A finite set D is ℓ -dissociated, where $\ell \geq 1$ is an integer, if every relation

$$\sum_{f \in D} a_f f = 0, \quad a_f \in \{-\ell, \dots, \ell\},$$

has all coefficients zero.

Proposition II.266.4 (A quantitative dissociated radius). *Let D be ℓ -dissociated with $q = |D| \geq 1$, and let $F = D \cup (-D)$. Put*

$$c_\ell = 2^{-\ell} \binom{2\ell}{\ell}, \quad \delta_{\ell,q} = \frac{1}{2\pi} \arccos \left(1 - \left(\frac{c_\ell}{q} \right)^{1/\ell} \right). \quad (\text{II.266.5})$$

One character separates F at radius $\delta_{\ell,q}$. For graph defects, the resulting weighted Freiman restriction retains $\delta_{\ell,q} w(B)/k$.

Proof. On the compact dual of a finitely generated subgroup containing D , set

$$R(\chi) = \prod_{f \in D} \frac{(1 - \cos(2\pi\chi(f)))^\ell}{c_\ell}.$$

The constant term of $(1 - \cos(2\pi x))^\ell$ is c_ℓ . To verify it, write $z = e^{2\pi i x}$ and use

$$(1 - \cos(2\pi x))^\ell = 2^{-\ell}(1 - z)^\ell(1 - z^{-1})^\ell;$$

the constant coefficient is $2^{-\ell} \sum_{j=0}^{\ell} \binom{\ell}{j}^2 = c_\ell$. Every frequency in the product has coefficients between $-\ell$ and ℓ on D . Dissociativity and character orthogonality imply that R , and every product obtained by deleting one factor, has integral one. This remains valid with torsion because the stated dissociativity excludes precisely the relevant zero-frequency aliases.

Let $\delta = \max_\chi \min_{f \in D} \|\chi(f)\|_{\mathbb{T}}$, whose maximum exists by compactness. At every character, some f has distance at most δ . Multiply that covering inequality by $R \geq 0$ and integrate. On the event $\|\chi(f)\|_{\mathbb{T}} \leq \delta$, its own factor is at most $(1 - \cos(2\pi\delta))^\ell / c_\ell$. Removing that factor and integrating the remaining product gives

$$1 \leq q \frac{(1 - \cos(2\pi\delta))^\ell}{c_\ell}.$$

Solving proves (II.266.5). The expression lies in the domain of $\arccos$: $c_\ell \leq 2^\ell$ and $q \geq 1$. Localization gives the last assertion. $\square$

For $\ell = 1$, this is

$$\delta_{1,q} = \frac{1}{2\pi} \arccos(1 - 1/q) \sim \frac{1}{\pi\sqrt{2q}}.$$

It substantially exceeds the elementary $1/(2q)$ radius from the infinite-order case when the full representative set is dissociated and q is large. The hypothesis on the entire set is essential: finding a large dissociated subset does not by itself separate the remaining defects. Adapting the argument to partial dissociativity is a distinct research problem.

II.267 Source 54's chapter: one quadratic phase at Proposition 8.1 [write]

Added in the tenth write, 7 October 2026.

The next eight sections are source 54's Sections 2–9, the report's first chapter on Gowers's Section 8. Source 54's introduction is Section F.59; its integration notes and questions are in Part V.

What is proved. For odd N , $f = \mathbf{1}_A - \delta$, any nonempty $P \subseteq \mathbb{Z}/N\mathbb{Z}$ of size T and the hypothesis $\sum_{h \in P} |D_h(\lambda h + \mu)|^2 \geq \beta N^2 T$ of Proposition 8.1, *one* polynomial $\psi(z) = 2^{-1}\lambda z^2 + rz$ satisfies $\sum_s |\sum_{z \in P+s} f(z)e(-\psi(z)/N)| \geq \beta NT/(B\sigma) \geq (27/4)\beta NT$, $\sigma = \delta(1 - \delta)$, $B = \max(\delta, 1 - \delta)$, where Gowers's translate-dependent phases give $1/\sqrt{2}$ (Theorem F.59.1). The source is an exact mixed weighted Fourier factorization with coefficient one, optimal (Theorem II.268.2); quadratic refinements extend it to every finite abelian group, and even order needs doubled roots (Proposition II.270.2, Example II.270.3). The offset energy $\mathcal{E}_\eta(g)$ is at most $c_m \sigma^2$ with $c_m = 1, 1/2, 1/3, 1/4$ by the order m of the offset, with all equality cases (Theorem II.271.3), and near-extremizers are within $K_m \sqrt{\sigma} \varepsilon^{1/2}$ ($m \leq 4$) or $3\sqrt{\sigma} \varepsilon^{1/4}$ ($m \geq 5$) of an extremizer, both exponents optimal (Theorems II.272.5, II.272.1, Proposition II.272.4); several intercepts are governed by the clique number (Theorem II.273.1,

through the Motzkin–Straus theorem, credited). The constant $27/4$ is the uniform consequence of the amplitude estimate, not an optimum over balanced indicators, and no lower bound on each translate is claimed, as source 54 says.

Relation to the formal statements (checked at HEAD fe7165a3c). `proposition_8_1` (Section `ns08_09.lean`, line 29) has the exact companion `proposition_8_1_holds` (Proofs08QuadraticBias.lean, line 1069), for odd N and a progression P ; Theorem F.59.1 implies it (note after the theorem) and is not formalized.

II.268 An exact mixed and weighted Fourier factorization

Source 54, delivered Section 2.

II.268.1 Conventions and definitions

All groups in the article are finite and nonempty. Write $n = |G|$ and $\mathbb{E}_x = n^{-1} \sum_{x \in G}$. For $u : G \rightarrow \mathbb{C}$,

$$\widehat{u}(\xi) = \mathbb{E}_x u(x) e(-\xi(x)), \quad u(x) = \sum_{\xi \in \widehat{G}} \widehat{u}(\xi) e(\xi(x)), \quad \sum_{\xi} |\widehat{u}(\xi)|^2 = \|u\|_2^2. \quad (\text{II.268.1})$$

The Fourier sums are not normalized a second time.

Let $u, v : G \rightarrow \mathbb{C}$, $\eta \in \widehat{G}$, and $w : G \rightarrow [0, \infty)$. Write $T = \sum_h w(h)$. Define

$$A(h) = A_\eta(u, v)(h) = \mathbb{E}_x u(x+h) \overline{v(x)} e(-\eta(x)), \quad (\text{II.268.2})$$

$$a_\xi = \widehat{u}(\xi) \overline{\widehat{v}(\xi - \eta)}, \quad \mathcal{D}_\eta(u, v) = \sum_{\xi} |a_\xi|, \quad (\text{II.268.3})$$

$$\mathcal{W}_{\xi, w} u(s) = \sum_h w(h) u(s+h) e(-\xi(h)), \quad (\text{II.268.4})$$

$$\mathcal{L}_{\xi, w}(u, v) = \mathbb{E}_s |v(s)| |\mathcal{W}_{\xi, w} u(s)|, \quad H = \max_{\xi} \mathcal{L}_{\xi, w}(u, v), \quad (\text{II.268.5})$$

$$\mathcal{E}_{\eta, w}(u, v) = \sum_h w(h) |A(h)|^2. \quad (\text{II.268.6})$$

In the auto-correlation case we abbreviate $u = v = g$.

Lemma II.268.1 (The two exact identities). *For every $h \in G$,*

$$A(h) = \sum_{\xi} a_{\xi} e(\xi(h)). \quad (\text{II.268.7})$$

If

$$K_{\xi} = \mathbb{E}_s \overline{v(s)} e(-\eta(s)) \mathcal{W}_{\xi, w} u(s), \quad (\text{II.268.8})$$

then

$$\mathcal{E}_{\eta, w}(u, v) = \sum_{\xi} \overline{a_{\xi}} K_{\xi}. \quad (\text{II.268.9})$$

The right side of (II.268.9) is real and nonnegative.

Proof. Expand $u(x+h)$ and $\overline{v(x)}$ using (II.268.1). Character orthogonality in x leaves exactly the pairs of frequencies $(\xi, \xi - \eta)$, proving (II.268.7). Also,

$$K_\xi = \sum_h w(h)A(h)e(-\xi(h)).$$

Substituting this expression and interchanging the finite sums gives

$$\sum_\xi \overline{a_\xi} K_\xi = \sum_h w(h)A(h) \overline{\sum_\xi a_\xi e(\xi(h))} = \sum_h w(h)|A(h)|^2.$$

There are no convergence or exchange-of-limits issues. $\square$

Theorem II.268.2 (Common-frequency extraction). *With the preceding notation,*

$$\mathcal{E}_{\eta,w}(u, v) \leq \mathcal{D}_\eta(u, v) \max_\xi \mathcal{L}_{\xi,w}(u, v), \quad \mathcal{D}_\eta(u, v) \leq \|u\|_2 \|v\|_2. \quad (\text{II.268.10})$$

The coefficient one in the first inequality is optimal. If $\mathcal{D}_\eta(u, v) > 0$, there is a single frequency ξ with

$$\mathbb{E}_s |v(s)| |\mathcal{W}_{\xi,w} u(s)| \geq \frac{\mathcal{E}_{\eta,w}(u, v)}{\mathcal{D}_\eta(u, v)}. \quad (\text{II.268.11})$$

Proof. By Lemma II.268.1, the triangle inequality, and $|K_\xi| \leq \mathcal{L}_{\xi,w}(u, v)$,

$$\mathcal{E}_{\eta,w}(u, v) \leq \sum_\xi |a_\xi| \mathcal{L}_{\xi,w}(u, v) \leq \mathcal{D}_\eta(u, v) H.$$

Cauchy–Schwarz in the dual group and Parseval give the second inequality in (II.268.10). A maximizer exists because the dual group is finite. To see optimality, take $u = v = ce(\xi_0(\cdot))$ with $c \neq 0$ and $\eta = 0$. Then $\mathcal{E} = |c|^4 T$, $\mathcal{D} = |c|^2$, and $H = |c|^2 T$. $\square$

Remark II.268.3 (Why the phase need not depend on the translate). The order of operations matters. We first sum the nonnegative quantities $|a_\xi| \mathcal{L}_{\xi,w}$ over frequencies, and only then choose a frequency. The average over s is already inside $\mathcal{L}_{\xi,w}$. Replacing this by a pointwise maximum in s would give the weaker conclusion of a different frequency on each translate.

II.268.2 What spectral overlap retains

For $u = v = g$ and $\sigma = \|g\|_2^2 > 0$,

$$\frac{\mathcal{D}_\eta(g)}{\sigma} = \sum_\xi \sqrt{p_\xi p_{\xi-\eta}}, \quad p_\xi = \frac{|\widehat{g}(\xi)|^2}{\sigma}. \quad (\text{II.268.12})$$

This is the affinity of a probability distribution and its translate. In particular,

$$1 - \frac{\mathcal{D}_\eta(g)}{\sigma} = \frac{1}{2} \sum_\xi (\sqrt{p_\xi} - \sqrt{p_{\xi-\eta}})^2. \quad (\text{II.268.13})$$

The inequality $\mathcal{D}_\eta(g) \leq \sigma$ is an equality precisely when $p_\xi = p_{\xi-\eta}$ for every ξ . It is strict whenever the power spectrum is not invariant under the offset. A small positive overlap improves the extraction denominator; a zero overlap means $A \equiv 0$ and hence zero input energy. There is no division by zero in any positive-energy application.

II.269 Quadratic untwisting and the strengthened proposition

Source 54, delivered Section 3.

II.269.1 The sign and normalization calculation

Return temporarily to $G = \mathbb{Z}/N\mathbb{Z}$ with N odd. Interpret $2^{-1}\lambda$ in $\mathbb{Z}/N\mathbb{Z}$, and put

$$q(x) = 2^{-1}\lambda x^2/N \pmod{1}, \quad g(x) = f(x)e(-q(x)), \quad \eta(x) = \mu x/N.$$

The defining property of q is

$$q(x+h) - q(x) - q(h) = \lambda hx/N \pmod{1}. \quad (\text{II.269.1})$$

In (F.59.2), set $x = y + h$ and use (II.269.1). The result is the exact identity

$$N^{-1}D_h(\lambda h + \mu) = e(-q(h) - \mu h/N)A_\eta(g)(h). \quad (\text{II.269.2})$$

This calculation is useful because the sign of h in the original backward derivative need not agree with a forward-difference convention. The harmless factor on the right has absolute value one.

Proof of Theorem F.59.1. Choose $w = \mathbf{1}_P$. By (II.269.2), the hypothesis says $\mathcal{E}_{\eta,w}(g) \geq \beta T$. Theorem II.268.2 provides one frequency r with

$$\mathbb{E}_s |g(s)| |\mathcal{W}_{r,w}g(s)| \geq \beta T / \mathcal{D}_\eta(g) \geq \beta T / \sigma.$$

Since $|g(s)| = |f(s)|$ and translation $h \mapsto z = s + h$ is a bijection,

$$|\mathcal{W}_{r,w}g(s)| = \left| \sum_{z \in P+s} f(z) e(-(2^{-1}\lambda z^2 + rz)/N) \right|.$$

Multiplication by N proves (F.59.6); then $|f(s)| \leq B$ proves the first inequality in (F.59.7). Finally let $t = \min(\delta, 1 - \delta)$. Then $B\sigma = t(1 - t)^2$ for $0 < t \leq 1/2$. Its derivative is $(1 - t)(1 - 3t)$, so its maximum is $4/27$, attained at $t = 1/3$. This proves the final inequality. $\square$

Remark II.269.1 (Sharp and non-sharp assertions). Coefficient one is sharp in the mixed extraction theorem. The factor $27/4$ is the exact uniform consequence of the displayed amplitude estimate, not a proved optimum over all balanced indicators. The constant-one extremizer used in Theorem II.268.2 is generally not a balanced indicator on an odd-order group. Those distinctions prevent an unwarranted sharpness claim.

The new unweighted coefficient tends to 8 as the density tends to $1/2$ (exact half-density is impossible for odd N). At density $1/3$ or $2/3$, it is $27/4$. The gain is accompanied by a strictly stronger common-polynomial conclusion, not merely by a numerical change. No lower bound on N is required beyond the stated algebraic hypotheses.

II.269.2 The form for arbitrary weights and bounded functions

Let f now be any nonzero complex function, with $B = \|f\|_\infty$, $\sigma = \|f\|_2^2$, and $a = \|f\|_1$. Assume $T > 0$ and $\mathcal{E}_{\eta,w}(g) > 0$, so that $\mathcal{D}_\eta(g) > 0$. The same calculation gives, for one common phase,

$$\mathbb{E}_s |f(s)| F(s) \geq \frac{\mathcal{E}_{\eta,w}(g)}{\mathcal{D}_\eta(g)}, \quad \mathbb{E}_s F(s) \geq \frac{\mathcal{E}_{\eta,w}(g)}{B \mathcal{D}_\eta(g)}, \quad (\text{II.269.3})$$

where

$$F(s) = \left| \sum_h w(h) f(s+h) e(-q(s+h) - \xi(s+h)) \right|.$$

In particular, at least one translate has

$$\frac{F(s)}{T} \geq \frac{\mathcal{E}_{\eta,w}(g)}{T a \mathcal{D}_\eta(g)}. \quad (\text{II.269.4})$$

For a balanced indicator, $a = 2\sigma$, so an energy lower bound βT yields $F(s)/T \geq \beta/(2\sigma^2)$. This is an existence statement for one translate, not an average bound of that size.

II.270 Every finite abelian group, including even order

Source 54, delivered Section 4.

Definition II.270.1. A quadratic refinement of a symmetric biadditive map $b : G \times G \rightarrow \mathbb{T}$ is a function $q : G \rightarrow \mathbb{T}$ with $q(0) = 0$ and

$$q(x+y) - q(x) - q(y) = b(x, y). \quad (\text{II.270.1})$$

Such a function is a quadratic phase polynomial in the additive-difference sense. It need not be represented by a polynomial with coefficients in $\mathbb{Z}/N\mathbb{Z}$, even when $G = \mathbb{Z}/N\mathbb{Z}$.

Proposition II.270.2 (Existence and the precise ambiguity). *Every symmetric biadditive $b : G \times G \rightarrow \mathbb{T}$ has a quadratic refinement. Any two refinements differ by a character of G .*

Proof. Write $G = \bigoplus_{i=1}^d \mathbb{Z}/n_i\mathbb{Z}$ with generators e_i , and let $b_{ij} = b(e_i, e_j)$. Choose $a_i \in \mathbb{T}$ such that

$$n_i a_i = -\binom{n_i}{2} b_{ii}. \quad (\text{II.270.2})$$

Such a choice exists because multiplication by n_i on $\mathbb{T}$ is surjective. For integer representatives x_i , set

$$q(x) = \sum_i \left(x_i a_i + \binom{x_i}{2} b_{ii} \right) + \sum_{i < j} x_i x_j b_{ij}. \quad (\text{II.270.3})$$

Changing x_i to $x_i + n_i$ changes this expression by

$$n_i a_i + \binom{n_i}{2} b_{ii} + n_i x_i b_{ii} + \sum_{j \neq i} n_i x_j b_{ij} = 0.$$

Here $n_i b_{ij} = 0$ follows from biadditivity. Thus q is well-defined on G . Expanding the binomial terms verifies (II.270.1). Conversely, subtracting (II.270.1) for two refinements proves that their difference is additive and hence a character. $\square$

For $f : G \rightarrow \mathbb{C}$, define the derivative coefficient selected by b and an intercept $\eta \in \widehat{G}$ by

$$C_{b,\eta}(h) = \mathbb{E}_x f(x+h) \overline{f(x)} e(-b(h,x) - \eta(x)). \quad (\text{II.270.4})$$

If q refines b and $g = f e(-q)$, then

$$C_{b,\eta}(h) = e(q(h)) A_\eta(g)(h). \quad (\text{II.270.5})$$

Consequently Theorem II.268.2 and (II.269.3) hold unchanged, with a single phase $q + \xi$ on every translate. Symmetry is the correct hypothesis: the polarization of any function satisfying (II.270.1) is necessarily symmetric. No assertion about nonsymmetric selectors is being hidden in the untwisting step.

II.270.1 The even cyclic case

On $\mathbb{Z}/N\mathbb{Z}$ with N even, the function

$$q(x) = \frac{\lambda x^2}{2N} \pmod{1} \quad (\text{II.270.6})$$

is well-defined, since $q(x+N) - q(x) = \lambda x + \lambda N/2$ is an integer. It refines $b(x,h) = \lambda x h/N$. When λ is odd, it generally takes values in $(2N)^{-1}\mathbb{Z}/\mathbb{Z}$ rather than in $N^{-1}\mathbb{Z}/\mathbb{Z}$.

Example II.270.3 (Why the target phase class must change). Let $G = \mathbb{Z}/2\mathbb{Z}$ and $f(0) = 1$, $f(1) = i$. Take $q(0) = 0$, $q(1) = 1/4$, so $f = e(q)$ and the untwisted function is identically one. For $b(x,h) = xh/2$ and $\eta = 0$, one has $|C_{b,0}(h)| = 1$ for both shifts. On the full group, extraction by q has normalized correlation one.

Every polynomial over $\mathbb{Z}/2\mathbb{Z}$, however, induces an affine function on its two points. Its exponential therefore has values proportional to $(1, 1)$ or $(1, -1)$. The largest normalized correlation with f is $|(1 \pm i)/2| = 1/\sqrt{2}$. Thus a constant-one strengthening for general bounded functions cannot retain the classical $\mathbb{Z}/2\mathbb{Z}$ polynomial phase class. The doubled-root refinement is essential.

Changing the refinement q changes it by a character and merely reindexes the maximizing frequency ξ . Thus the available family of common phases is intrinsic to b .

II.271 Exact global offset ceilings and equality cases

Source 54, delivered Section 5.

Lemma II.271.1 (Power-spectrum identity). *For every $g : G \rightarrow \mathbb{C}$ and $\eta \in \widehat{G}$,*

$$\mathcal{E}_\eta(g) = \sum_\xi |\widehat{g}(\xi)|^2 |\widehat{g}(\xi - \eta)|^2. \quad (\text{II.271.1})$$

In particular $\mathcal{E}_\eta(g) = \mathcal{E}_{-\eta}(g)$, and

$$\sum_{\eta \in \widehat{G}} \mathcal{E}_\eta(g) = \|g\|_2^4. \quad (\text{II.271.2})$$

Proof. Apply Parseval in h to (II.268.7) with $u = v = g$. The symmetry follows by reindexing ξ . Summing over all η factors the resulting double sum into two copies of $\sum_\xi |\widehat{g}(\xi)|^2$. $\square$

[write] **7 October 2026; sources 54–59.** *Relation to source 31.* The function $(h, \eta) \mapsto |A_\eta(g)(h)|^2/\sigma^2$ is source 31's phase-space function q_f (with $f = g/\sqrt{\sigma}$; Section II.104), and (II.271.1) is its sum over h at the fixed frequency η ; summing also over η gives (II.271.2) and source 31's total-mass identity in Lemma II.105.1. Source 54 does not cite source 31 (note in Section F.59).

We will need an explicit triangle-free graph certificate. Besides proving the ceiling, it will select the four frequencies from which the stability theorem is obtained.

Lemma II.271.2 (A weighted triangle-free certificate). *Let Γ be a finite simple triangle-free graph. Assign probabilities $p_v \geq 0$ to its vertices, with $\sum_v p_v = 1$, and put*

$$F = \sum_{\{u,v\} \in E(\Gamma)} p_u p_v, \quad d_v = \sum_{u: \{u,v\} \in E(\Gamma)} p_u.$$

Then $F \leq 1/4$. More precisely,

$$F - 4F^2 = \sum_{\{u,v\} \in E} p_u p_v (1 - d_u - d_v) + \left(\sum_v p_v d_v^2 - (2F)^2 \right), \quad (\text{II.271.3})$$

and both terms on the right are nonnegative. If $F > 0$, some edge with $p_u p_v > 0$ satisfies

$$1 - d_u - d_v \leq 1 - 4F. \quad (\text{II.271.4})$$

Proof. The neighborhoods of the endpoints of an edge are disjoint because there are no triangles. Therefore $d_u + d_v \leq 1$ on every edge. Furthermore,

$$\sum_v p_v d_v = 2F, \quad \sum_{\{u,v\} \in E} p_u p_v (d_u + d_v) = \sum_v p_v d_v^2.$$

These equations give (II.271.3); the second term is a variance. Hence $F - 4F^2 \geq 0$ and $F \leq 1/4$. The average of $1 - d_u - d_v$ over edges with weights $p_u p_v / F$ is at most $1 - 4F$, which proves (II.271.4). $\square$

Theorem II.271.3 (The exact extremizers). *Let $g \neq 0$, $\sigma = \|g\|_2^2$, $m = \text{ord}(\eta)$, and $p_\xi = |\hat{g}(\xi)|^2/\sigma$. Then $\mathcal{E}_\eta(g) \leq c_m \sigma^2$, with c_m as in (F.59.10). Equality is characterized by the following table. In every row all unlisted frequencies have zero power, and the phases of nonzero Fourier coefficients are arbitrary.*

m	c_m	Power-spectrum description
1	1	A single frequency has mass 1.
2	1/2	A single η -orbit has masses 1/2, 1/2.
3	1/3	A single η -orbit has masses 1/3, 1/3, 1/3.
4	1/4	A single four-cycle carries all the mass; each of its two alternating pairs has total mass 1/2.
≥ 5	1/4	For some ξ_0 , the frequencies $\xi_0 - \eta, \xi_0, \xi_0 + \eta$ carry masses $a, 1/2, 1/2 - a$, where $0 \leq a \leq 1/2$.

Proof. Translation by η partitions $\hat{G}$ into cycles of length m . By Lemma II.271.1, we must maximize $\sum_\xi p_\xi p_{\xi-\eta}$.

If $m = 1$, this is $\sum p_\xi^2 \leq 1$, with equality precisely for a point mass. If $m = 2$, a cycle with masses a, b contributes $2ab$. Writing $t = a + b$, one has $2ab \leq t^2/2$. Summing over cycles gives at most $1/2$; equality requires one cycle and $a = b = 1/2$.

For $m = 3$, a cycle with masses a, b, c contributes $ab + bc + ca$. This is at most $(a + b + c)^2/3$, with equality when $a = b = c$. Again all mass must belong to one cycle to attain $1/3$.

For $m \geq 4$, the graph of these cycles is simple and triangle-free, and each undirected edge appears once in the offset sum. Lemma II.271.2 gives the upper bound $1/4$. If equality holds, choose an edge as in (II.271.4). Its two neighborhoods carry all the mass. When $m = 4$, their union is one four-cycle. If its alternating-pair masses are U, V , the energy is UV , so $U = V = 1/2$.

When $m \geq 5$, the union consists of four consecutive vertices, inducing a path. Let their masses in path order be a, b, c, d . Writing $U = a + c$ and $V = b + d$ gives

$$ab + bc + cd = UV - ad \leq \frac{1}{4}. \quad (\text{II.271.5})$$

Equality forces $U = V = 1/2$ and $ad = 0$. Removing the zero endpoint leaves a three-vertex star with center mass $1/2$. This gives the last row of the table. Conversely every distribution in the table attains the claimed bound by direct substitution. $\square$

Remark II.271.4. For general graphs, broad descriptions of every Motzkin–Straus maximizer require care; erroneous classifications in earlier literature have been corrected [225]. The proof above does not appeal to such a classification. It derives the equality cases directly from the special cycle geometry.

The constants remain optimal when one restricts to functions with $\|g\|_\infty \leq 1$, provided the L^2 norm is not prescribed in advance: construct a Fourier polynomial from a listed spectrum and scale it by a sufficiently small nonzero constant. The ratio $\mathcal{E}_\eta(g)/\|g\|_2^4$ is unchanged. This does not prove optimality for balanced indicators of a prescribed density.

II.272 Dimension-free stability and the fourth-root obstruction

Source 54, delivered Section 6.

II.272.1 The cycle stability lemma

Theorem II.272.1 (Recovery of a three-frequency star). *Let Γ be a disjoint union of cycles of a common length $m \geq 5$. Let p be a probability distribution on its vertices and suppose*

$$F = \sum_{\{u,v\} \in E} p_u p_v \geq \frac{1}{4} - \varepsilon, \quad 0 \leq \varepsilon \leq \frac{1}{16}.$$

There is an extremal star distribution q from the last row of Theorem II.271.3 for which

$$\|p - q\|_1 \leq 8\varepsilon + 4\sqrt{3\varepsilon}. \quad (\text{II.272.1})$$

The exponent $1/2$ in this probability-distance estimate is optimal.

Proof. Since $F \geq 3/16 > 0$, choose an edge using Lemma II.271.2. The union of the neighborhoods of its endpoints is a four-vertex path S . Its outside mass r satisfies

$$r \leq 1 - 4F \leq 4\varepsilon. \quad (\text{II.272.2})$$

Write the masses on the path as a, b, c, d in order; their sum is $t = 1 - r$. Put $U = a + c$, $V = b + d$, and $D = |U - V|$.

The contribution of edges crossing from S to its complement is at most rt , simply by allowing every possible pair. By the triangle-free bound applied to the outside distribution of total mass r , its internal contribution is at most $r^2/4$. Consequently, using (II.271.5),

$$\begin{aligned} ad + \frac{D^2}{4} &= \frac{t^2}{4} - F + F_{\text{outside}} \\ &\leq \frac{(1-r)^2}{4} - \left(\frac{1}{4} - \varepsilon\right) + r(1-r) + \frac{r^2}{4} \\ &= \varepsilon + \frac{r}{2} - \frac{r^2}{2} \leq 3\varepsilon. \end{aligned} \quad (\text{II.272.3})$$

Here F_{outside} includes both crossing and wholly outside edges. It follows that

$$\min(a, d) \leq \sqrt{3\varepsilon}, \quad D \leq 2\sqrt{3\varepsilon}. \quad (\text{II.272.4})$$

Suppose $a \leq d$; the other case is symmetric. Delete the outside mass and the left endpoint. The retained star has center mass c and total leaf mass $b + d$. Define q to have center mass $1/2$ and distribute the remaining $1/2$ between the two leaves in the same proportion as their old masses. If both old leaf masses vanish, choose either leaf; the following distance calculation still holds. Let $R = r + a$ be the total deleted mass. Then

$$\begin{aligned} \|p - q\|_1 &= R + |c - 1/2| + |b + d - 1/2| \\ &= R + \max\{R, |c - b - d|\} \\ &\leq R + \max\{R, D + a\} \\ &\leq 2R + D \\ &\leq 8\varepsilon + 4\sqrt{3\varepsilon}. \end{aligned} \quad (\text{II.272.5})$$

The penultimate inequality uses $a \leq R$. The optimality assertion is proved in Proposition II.272.4. $\square$

II.272.2 From powers to functions

Lemma II.272.2 (Phase-preserving lifting). *Let $g \neq 0$, $\sigma = \|g\|_2^2$, and $p_\xi = |\widehat{g}(\xi)|^2/\sigma$. For any probability distribution q on $\widehat{G}$, there is g_q with Fourier powers σq for which*

$$\|g - g_q\|_2^2 = \sigma \sum_{\xi} (\sqrt{p_\xi} - \sqrt{q_\xi})^2 \leq \sigma \|p - q\|_1. \quad (\text{II.272.6})$$

Proof. Use the phase of $\widehat{g}(\xi)$ for the new coefficient $\widehat{g}_q(\xi) = \sqrt{\sigma q_\xi}$ times that phase. If the old coefficient vanishes, choose an arbitrary phase. Parseval gives the identity. The inequality follows pointwise from $(\sqrt{a} - \sqrt{b})^2 \leq |a - b|$ for $a, b \geq 0$. $\square$

Corollary II.272.3 (Fourth-root recovery). *If $m \geq 5$ and $\mathcal{E}_\eta(g) \geq (1/4 - \varepsilon)\sigma^2$ with $0 \leq \varepsilon \leq 1/16$, an exact extremizer g_* satisfies*

$$\|g - g_*\|_2 \leq \sqrt{\sigma} \sqrt{8\varepsilon + 4\sqrt{3}\varepsilon} \leq 3\sqrt{\sigma} \varepsilon^{1/4}. \quad (\text{II.272.7})$$

Proof. Combine Lemma II.271.1, Theorem II.272.1, and Lemma II.272.2. For the last inequality use $8\varepsilon \leq 2\sqrt{\varepsilon}$ and $2 + 4\sqrt{3} < 9$. $\square$

The recovered function has the same L^2 norm as g . The proof does not preserve an additional pointwise bound, a real-valued condition, or the property of being a balanced indicator. Such preservation is a separate problem and is not asserted here.

Proposition II.272.4 (The sharp obstruction). *Fix any offset order $m \geq 5$. For every $0 < t \leq 1/4$, there is a function g_t with $\|g_t\|_2^2 = \sigma$ such that*

$$\mathcal{E}_\eta(g_t) = \left(\frac{1}{4} - t^2\right) \sigma^2, \quad \inf_{g_*} \|g_t - g_*\|_2 \geq \sqrt{\sigma} t^{1/2}, \quad (\text{II.272.8})$$

where the infimum is over all exact extremizers of the same norm. The corresponding power distribution is at ℓ^1 distance at least $2t$ from every extremal power distribution. These examples may all be chosen with $\|g_t\|_\infty \leq 1$ by fixing $\sigma = 1/4$.

Proof. Place the four consecutive Fourier powers

$$(t, 1/2 - t, 1/2 - t, t) \quad (\text{II.272.9})$$

on one η -orbit and put zero elsewhere. The induced support is a path, so its edge sum is

$$t(1/2 - t) + (1/2 - t)^2 + t(1/2 - t) = 1/4 - t^2.$$

Every extremal star misses at least mass t of this distribution: a three-vertex star cannot contain all four consecutive vertices, and every positive mass is at least t . A star centered at either of the two internal vertices misses exactly one endpoint, so t is the exact minimum omitted mass.

For functions, the omitted Fourier coefficients alone contribute at least σt to squared L^2 distance, independent of their phases and of any changes on the retained frequencies. For power probabilities, deleting mass at least t and restoring total mass one costs at least $2t$ in ℓ^1 . Finally, choose Fourier coefficients of magnitudes $\sqrt{\sigma p_\xi}$. Their absolute sum is at most $2\sqrt{\sigma}$, by Cauchy–Schwarz on four frequencies. Thus $\sigma = 1/4$ ensures the desired pointwise bound. $\square$

With $\varepsilon = t^2$, the lower bound is exactly of fourth-root order. In particular, an $O(\varepsilon^{1/2})$ theorem for all nonzero offsets is false, even for uniformly bounded functions on a fixed cyclic group of order at least five.

II.272.3 Orders at most four

Theorem II.272.5 (Square-root recovery for short orbits). *For $m \in \{1, 2, 3, 4\}$, under the hypotheses of Theorem F.59.2, there is an exact extremizer of norm $\sqrt{\sigma}$ satisfying the first line of (F.59.11).*

Proof. Work with the power distribution p and lift the construction by preserving Fourier phases. All squared distances below are divided by σ .

Order one. Since $\sum p_\xi^2 \geq 1 - \varepsilon$, some $p_\xi \geq 1 - \varepsilon$. Lifting the point mass at that frequency gives squared distance $2 - 2\sqrt{p_\xi} \leq 2\varepsilon$.

Order two. Let t_j denote orbit masses. The energy is at most $\frac{1}{2} \sum_j t_j^2$, hence a largest orbit has mass at least $1 - 2\varepsilon$. Let $r \leq 2\varepsilon$ be the outside mass and a, b the masses of the chosen pair. Outside orbits contribute at most $r^2/2$, so

$$(a - b)^2 = (1 - r)^2 - 4ab \leq 2\varepsilon - 2r + 2r^2 \leq 2\varepsilon.$$

Replace a, b by $1/2, 1/2$. Using $(\sqrt{x} - \sqrt{y})^2 \leq (x - y)^2/y$ for $y > 0$, the inside squared Hellinger distance is at most

$$2((a - 1/2)^2 + (b - 1/2)^2) = r^2 + (a - b)^2.$$

Including the outside mass gives at most $4\varepsilon + 4\varepsilon^2 \leq 5\varepsilon$.

Order three. A largest orbit has mass $t \geq 1 - 3\varepsilon$; put $r = 1 - t$. For its three masses a_i , let $V = \sum_{i=1}^3 (a_i - t/3)^2$. Its contribution is $t^2/3 - V/2$, while the other orbits contribute at most $r^2/3$. The energy assumption gives $V \leq 2\varepsilon$. Replacing the three masses by $1/3$ has squared Hellinger cost at most

$$r + 3 \sum_i (a_i - 1/3)^2 = r + 3V + r^2 \leq 9\varepsilon + 9\varepsilon^2 \leq 10\varepsilon.$$

Order four. Lemma II.271.2 selects one four-cycle with outside mass $r \leq 4\varepsilon$. Let U, V be its alternating-pair masses. The outside energy is at most $r^2/4$, whence

$$(U - V)^2 \leq 4\varepsilon - 2r + 2r^2 \leq 4\varepsilon.$$

Scale the two alternating-pair distributions to total mass $1/2$ each, preserving their internal proportions. The inside squared Hellinger cost is $(\sqrt{U} - 1/\sqrt{2})^2 + (\sqrt{V} - 1/\sqrt{2})^2$, at most $r^2 + (U - V)^2$. With the outside mass, the total is at most $8\varepsilon + 16\varepsilon^2 \leq 9\varepsilon$. These estimates give the four constants claimed. $\square$

Proposition II.272.6 (Optimality of the short-orbit exponents). *The exponent $1/2$ in Theorem II.272.5 cannot be increased uniformly. For $m = 1$ assume $|G| \geq 2$.*

Proof. For $m = 1$, the distribution $(1 - t, t)$ has deficit $2t - 2t^2$ and every extremal point mass omits at least t , giving L^2 distance at least $\sqrt{\sigma t}$. For $m = 2$, the pair $(1/2 + t, 1/2 - t)$ has deficit $2t^2$. The unique extremal powers on that orbit are $(1/2, 1/2)$, and the Hellinger distance is bounded below by a positive constant times $|t|$ for small t . For $m = 3$, use $(1/3 + t, 1/3 - t, 1/3)$; the deficit is t^2 and the same conclusion follows. For $m = 4$, put masses $1/2 + t, 1/2 - t$ on one edge. The energy deficit is t^2 , and every extremizer must have total mass $1/2$ on each alternating pair. Cauchy–Schwarz on each pair bounds the Hellinger distance below by that of the pair totals, again a positive constant times $|t|$. Fourier phases can be arbitrary, since choosing matching phases minimizes distance. $\square$

II.273 Several intercepts and the classical clique parameter

Source 54, delivered Section 7.

The single-offset problem has a particularly simple graph, but the spectral identity accommodates more than one intercept without additional analytic losses.

Theorem II.273.1 (Simultaneous offset energy). *Let $S \subseteq \widehat{G} \setminus \{0\}$ satisfy $S = -S$. Let Γ_S be the simple graph on $\widehat{G}$ in which ξ, ζ are adjacent exactly when $\xi - \zeta \in S$. If its clique number is ω , then*

$$\sum_{\eta \in S} \mathcal{E}_\eta(g) \leq \left(1 - \frac{1}{\omega}\right) \|g\|_2^4. \quad (\text{II.273.1})$$

The constant is best possible. For $S = \emptyset$, take $\omega = 1$.

Proof. For $g \neq 0$, Lemma II.271.1 writes the left side divided by σ^2 as

$$2 \sum_{\{\xi, \zeta\} \in E(\Gamma_S)} p_\xi p_\zeta.$$

The classical Motzkin–Straus theorem [224] says that the maximum of this expression over probability distributions is $1 - 1/\omega$. For completeness, its extremal-value argument is short. If two nonadjacent vertices both have positive mass, redistribute their combined mass entirely to whichever has the larger weighted neighbor sum. The objective does not decrease, because it is linear in this redistribution and there is no edge between the two vertices. Repeating reduces the support to a clique. On a clique of size $k \leq \omega$, the objective is $1 - \sum p_\xi^2 \leq 1 - 1/k \leq 1 - 1/\omega$. Uniform mass on a maximum clique attains equality, and it can be realized as the normalized Fourier power spectrum of a function. $\square$

Taking $S = \{\eta, -\eta\}$ recovers the nonzero-offset values in (F.59.10), remembering that this set has one element if $\text{ord}(\eta) = 2$. The equality classification for general S is not inferred from the support-reduction proof: existence of a clique-supported maximizer does not imply that every maximizer has clique support. The four-cycle row of Theorem II.271.3 already demonstrates the distinction.

Example II.273.2 (A subgroup of intercepts). If $H \leq \widehat{G}$ and $S = H \setminus \{0\}$, then Γ_S is a disjoint union of complete graphs on the cosets of H . Equation (II.273.1) has constant $1 - 1/|H|$. For $|H| > 1$, equality means that the power spectrum is uniform on one coset of H . Indeed, a coset of mass t contributes at most $(1 - 1/|H|)t^2$, with equality at equal masses, and summing forces all the mass into one coset.

More generally, nonnegative symmetric weights $c_\eta = c_{-\eta}$ give the exact identity

$$\sum_{\eta} c_\eta \mathcal{E}_\eta(g) = \sigma^2 \sum_{\xi, \zeta} c_{\xi - \zeta} p_\xi p_\zeta. \quad (\text{II.273.2})$$

This turns any valid upper certificate for the finite quadratic program on the right into a derivative-spectrum bound. No claim is made that its value is always determined by an unweighted clique number.

II.274 Localization, robustness, and many useful translates

Source 54, delivered Section 8.

II.274.1 How a local hypothesis forces global structure

Let $P \subseteq G$ have density $\rho = |P|/|G| > 0$, and assume

$$\frac{1}{|P|} \sum_{h \in P} |A_\eta(g)(h)|^2 \geq \beta. \quad (\text{II.274.1})$$

Positivity and the global ceiling give

$$\rho\beta \leq \mathcal{E}_\eta(g) \leq c_m \sigma^2, \quad \beta \leq \sigma^2 \min\{1, c_m/\rho\}. \quad (\text{II.274.2})$$

The second bound also uses $|A_\eta(g)(h)| \leq \sigma$. Thus, if $\rho\beta \geq (c_m - \varepsilon)\sigma^2$, the entire function is close to the extremal Fourier family specified in Theorem F.59.2. This gives an explicit local-to-global recovery criterion, with the same stability exponents.

Corollary II.274.1 (Certification that the intercept is zero). *In (II.274.1), the intercept must be zero if $\rho\beta > \sigma^2/2$. If $|G|$ is odd, it suffices that $\rho\beta > \sigma^2/3$. If $|G|$ is odd and $3 \nmid |G|$, it suffices that $\rho\beta > \sigma^2/4$.*

Proof. Every nonzero character has order at least two; on an odd-order group it has odd order at least three; and in the last case its order is at least five. Apply (F.59.10). $\square$

On $\mathbb{Z}/p\mathbb{Z}$ with $p \geq 5$ prime, this applies directly to the intercept μ in the quadratic selector after untwisting. The thresholds in the corollary are sharp as global L^2 -normalized statements when the corresponding character orders occur. They are not claimed sharp for every prescribed proper subset P .

Example II.274.2 (The density factor cannot be discarded). Take $g = \mathbf{1}_{\{0\}}$ on any group, let $P = \{0\}$, and choose any nonzero character η . Then $\sigma = 1/|G|$ and $A_\eta(g)(0) = \sigma$. Thus the local ratio is $\beta/\sigma^2 = 1$, despite the nonzero intercept. The global energy is only $\sigma^2/|G|$. This bounded example rules out applying the nonzero-offset ceiling directly to an arbitrary local average without its coverage factor ρ .

For a weight $0 \leq w \leq L$, the analogous bound is $\mathcal{E}_{\eta,w}(g) \leq L|G|c_m\sigma^2$. Its effective coverage is $T/(L|G|)$. This observation handles repeated or weighted shifts without treating multiplicity as additional independent coverage.

II.274.2 A complementary global correlation bound

The local extraction does not lose a factor ρ , because it keeps the correlation localized. A different conclusion follows by paying that factor.

Proposition II.274.3. *Under (II.274.1), some frequency ξ satisfies*

$$|\widehat{g}(\xi)| \geq \sqrt{\rho\beta/\sigma}. \quad (\text{II.274.3})$$

For a quadratic untwisting $g = fe(-q)$, this is a global correlation of f with the single phase $q + \xi$.

Proof. By (II.271.1), $\mathcal{E}_\eta(g) \leq \max_\xi |\widehat{g}(\xi)|^2 \sum_\zeta |\widehat{g}(\zeta - \eta)|^2 = \sigma \max_\xi |\widehat{g}(\xi)|^2$. Now use $\mathcal{E}_\eta(g) \geq \rho\beta$. $\square$

II.274.3 Noisy derivative data

Proposition II.274.4 (An L^2 -robust input). *Let $T = \sum w > 0$. Suppose observations $\tilde{A}(h)$ satisfy*

$$\left(\frac{1}{T} \sum_h w(h) |\tilde{A}(h)|^2 \right)^{1/2} \geq \sqrt{\beta}, \quad \left(\frac{1}{T} \sum_h w(h) |\tilde{A}(h) - A_\eta(g)(h)|^2 \right)^{1/2} \leq \zeta.$$

Then every extraction conclusion holds with the effective energy parameter (positive-energy denominators are used only when this is positive)

$$\beta_{\text{eff}} = (\sqrt{\beta} - \zeta)_+^2. \quad (\text{II.274.4})$$

Proof. The reverse triangle inequality in the weighted finite L^2 space gives $\|A_\eta(g)\|_{L^2(w/T)} \geq (\sqrt{\beta} - \zeta)_+$. Squaring gives the required energy bound. $\square$

This statement is about error in complex correlation data. It does not assert that a small numerical error in a frequency label always has a small effect; that requires a separate estimate translating label error into correlation error. Likewise, an approximately symmetric bilinear selector does not automatically have an exact quadratic refinement.

For an erasure variant, if an observed collection has total weight T and the unusable shifts have weight at most θT , then $|A_\eta(g)(h)|^2 \leq \sigma^2$ shows that deleting them loses at most $\theta\sigma^2 T$ of true energy. Hence the effective lower bound on the retained energy, still normalized by the original T , is $(\beta - \theta\sigma^2)_+ T$.

II.274.4 Quantitative abundance for the chosen common phase

Use the phase selected in (II.269.3), put $Z(s) = F(s)/T$, and suppose $\mathcal{E}_{\eta,w}(g) \geq \beta T$. Then $0 \leq Z \leq B$ and $\mathbb{E}_s |f(s)| Z(s) \geq \beta / \mathcal{D}_\eta(g)$. For $0 \leq \tau < B$, let $S_\tau = \{s : Z(s) > \tau\}$. Splitting the last average over S_τ and its complement gives

$$\mathbb{E}_s |f(s)| \mathbf{1}_{S_\tau}(s) \geq \frac{(\beta / \mathcal{D}_\eta(g) - \tau \|f\|_1)_+}{B - \tau}, \quad \frac{|S_\tau|}{|G|} \geq \frac{(\beta / \mathcal{D}_\eta(g) - \tau \|f\|_1)_+}{B(B - \tau)}. \quad (\text{II.274.5})$$

This is an elementary capped-mean consequence, not a new general selection inequality. Its useful feature here is that all selected translates use the *same* quadratic phase. For a balanced indicator one may substitute $\|f\|_1 = 2\sigma$ and $\mathcal{D}_\eta(g) \leq \sigma$.

II.275 An exact deficit and a constructive selector

Source 54, delivered Section 9.

II.275.1 Where near equality can lose mass

Keep the notation of Section II.268. Since $|K_\xi| \leq \mathcal{L}_\xi \leq H$, the factorization yields the exact nonnegative decomposition

$$\mathcal{D}H - \mathcal{E} = \sum_\xi |a_\xi| (H - \mathcal{L}_\xi) + \sum_\xi (|a_\xi| \mathcal{L}_\xi - \Re(\overline{a_\xi} K_\xi)). \quad (\text{II.275.1})$$

There are two different losses. The first measures concentration on frequencies with poor average window score. The second is phase cancellation inside the corresponding correlations.

For $a_\xi \neq 0$, set $\theta_\xi = a_\xi/|a_\xi|$ and $Z_\xi(s) = \overline{v(s)}e(-\eta(s))\mathcal{W}_{\xi,w}u(s)$. The second summand has the more explicit form

$$|a_\xi|\mathbb{E}_s \left(|Z_\xi(s)| - \Re(\overline{\theta_\xi}Z_\xi(s)) \right). \quad (\text{II.275.2})$$

Thus exact equality forces the nonzero $Z_\xi(s)$ to have the required common argument whenever the corresponding spectral weight is positive, in addition to maximal average window score. Equation (II.275.1) is also a certificate for numerical or formal implementations: no estimate is hidden between its two sides.

If $\mathcal{D} > 0$, define $\pi_\xi = |a_\xi|/\mathcal{D}$. A random frequency sampled from π has expected score at least $\mathcal{E}/\mathcal{D}$. If $H > 0$ and $\mathcal{E} \geq (1 - \varepsilon)\mathcal{D}H$, then for every $0 < \tau < 1$,

$$\pi\{\xi : \mathcal{L}_\xi < (1 - \tau)H\} \leq \varepsilon/\tau. \quad (\text{II.275.3})$$

This is Markov's inequality applied to the first deficit term. It is a useful robustness statement about extraction, separate from the offset-energy stability theorem.

II.275.2 Finite implementation

A deterministic selector can be implemented as follows: compute $\widehat{u}$ and $\widehat{v}$; form the products a_ξ ; for each frequency compute $\mathcal{W}_{\xi,w}u$ and its score $\mathcal{L}_\xi$; retain a maximizing frequency and report the values in (II.268.9) and (II.275.1). The group Fourier transform and the chosen group representation are explicit input.

For $\mathbb{Z}/N\mathbb{Z}$, each window array is a cyclic correlation and can be computed by FFT. Scanning all N frequencies uses $O(N^2 \log N)$ arithmetic operations and $O(N)$ working memory when arrays are streamed. This is an arithmetic-operation bound, not an exact bit-complexity theorem for arbitrary complex inputs. With exact cyclotomic data, comparisons of scores require certified algebraic or interval arithmetic. The included implementation uses floating point for Fourier regression tests and exact integer arithmetic only for the graph certificates.

The stability selector for $m \geq 5$ is particularly simple. Compute neighbor masses d_v , choose an edge minimizing $1 - d_u - d_v$ among edges of positive weight, delete the smaller endpoint of its four-vertex neighborhood path, and rebalance the resulting star. All structural choices can be made from the power spectrum alone; Fourier phases enter only when reconstructing the nearby function.

II.276 Source 58's chapter: Lemma 9.3 on every cyclic group [write]

Added in the tenth write, 7 October 2026.

The next ten sections are source 58's Sections 2–11, on the restriction step of Gowers's Lemma 9.3. *Logically they continue the chapters of sources 17, 29 and 48* (Sections II.57, II.119 and II.186) and source 40's barrier (Section II.164); they are printed here because no number printed earlier may move. Source 58's introduction is Section F.63.

What is proved. For every $N \geq 1$, $B \subseteq \mathbb{Z}/N\mathbb{Z}$ with $b = |B| \geq (4/(\alpha\eta))^{54}$ and any map $\phi : B \rightarrow H$ to a finite abelian group with $\mathsf{G}_{16}(B, \phi) \geq \alpha b^{15}$, some $B' \subseteq B$ has $\mathsf{G}_{16}(B', \phi) \geq (1 - \eta)\mathsf{T}_{16}(B')$

and $T_{16}(B') \geq (\alpha\eta/4)^{54}b^{15}$ (Theorem A); Gowers's exponent is 2^{19} . It comes from an expectation theorem for amplified positive Fourier filters with explicit nonprincipal error $\tau_G(2D)(R_m^k - P^k)b^{m-2}$ and collision error $\binom{m}{2}R_2^kb^{m-2}$ (Theorem II.282.1), giving exponent $O(m \log m)$ for every balanced relation of length m on every cyclic group (Theorem II.282.3), arbitrary targets (Corollary II.284.2) and the quadruple input of Corollary 9.4 (Corollary II.286.2). Within normalized nonnegative filters the leading principal cost is exactly $m \log m / (2 \log a)$ when a discrepancy may have order a (Theorem B), the multi-layer counterpart of source 40's barrier. This answers source 17's composite-modulus question for one balanced relation (item II.8). The exponent 54 is a certified convenience, not optimal, and the prime-field exponents of sources 04 and 17 remain better there, as source 58 says.

Relation to the formal statements (checked at HEAD fe7165a3c). `lemma_9_3` and `corollary_9_4` (Sections08_09.lean, lines 77 and 91) have the exact companions `lemma_9_3_holds` and `corollary_9_4_holds` (Proofs09Restriction.lean, lines 376 and 416), on all moduli with an existential threshold and exponent 2^{19} ; source 58's counts are `additiveTupleCount`, `phiAdditiveTupleCount` and `GammaHomOfOrder` (Definitions.lean, lines 322, 327 and 334). Theorem A implies `lemma_9_3` with $N_0 = \lceil \beta^{-1}(4/(\alpha\eta))^{54} \rceil$ (Section II.286); it is not formalized.

II.277 Counts, normalizations, and finite Fourier analysis

Source 58, delivered Section 2.

Fix $m = 2r \geq 4$ and put

$$\varepsilon = (\underbrace{1, \dots, 1}_r, \underbrace{-1, \dots, -1}_r).$$

Groups are written additively. For $B \subseteq G$ and $\phi : B \rightarrow H$, define

$$\mathcal{T}_m(B) = \{\mathbf{x} \in B^m : \sum_{i=1}^m \varepsilon_i x_i = 0\}, \quad (\text{II.277.1})$$

$$\mathcal{G}_m(B, \phi) = \{\mathbf{x} \in \mathcal{T}_m(B) : \sum_{i=1}^m \varepsilon_i \phi(x_i) = 0\}. \quad (\text{II.277.2})$$

Their cardinalities are T_m and G_m . The remaining tuples form $\mathcal{U}_m$, with cardinality U_m . We call a tuple in $\mathcal{G}_m$ *respected* and one in $\mathcal{U}_m$ *bad*. The target discrepancy of a tuple is

$$\delta(\mathbf{x}) = \sum_i \varepsilon_i \phi(x_i) \in H. \quad (\text{II.277.3})$$

No injectivity hypothesis is placed on ϕ .

Lemma II.277.1 (Free-variable and collision counts). *For every finite abelian G and $B \subseteq G$ of size b ,*

$$T_m(B) \leq b^{m-1}. \quad (\text{II.277.4})$$

The number of additive tuples with at least one repeated entry is at most

$$\binom{m}{2} b^{m-2}. \quad (\text{II.277.5})$$

Every bad tuple has at least two distinct entries.

Proof. The first $m - 1$ entries determine the last, since its coefficient is a unit ± 1 . For the second assertion, fix a pair of positions $i < j$ that will be equal, and choose a position h outside that pair. Choose the $m - 2$ entries in all positions except j, h ; then $x_j = x_i$, and the coefficient $\varepsilon_h = \pm 1$ determines x_h . At most b^{m-2} tuples arise for this pair. Sum over the pairs, allowing overcounting. Finally, if every entry equals x , both balanced sums of its image equal $r\phi(x)$, so the tuple is respected. $\square$

The group of characters $G \rightarrow \mathbb{T} = \mathbb{R}/\mathbb{Z}$ is denoted $\widehat{G}$. With $e(t) = \exp(2\pi it)$, character orthogonality is

$$\mathbb{E}_{\chi \in \widehat{G}} e(\chi(x)) = \mathbf{1}_{x=0}. \quad (\text{II.277.6})$$

Here and below uniform averages on a finite set use total mass one. For completeness, if $x \neq 0$, the finite abelian decomposition supplies a character χ_0 with $e(\chi_0(x)) \neq 1$; translating the sum by χ_0 multiplies it by that number, forcing it to vanish.

On $\mathbb{T}$ integration always means normalized Haar integration. For functions on a finite abelian group V , the unnormalized Fourier transform is

$$\widehat{f}(\chi) = \sum_{x \in V} f(x) e(-\chi(x)).$$

This convention will be used only for the moment calculation in Section II.286; all random-character averages are normalized as in (II.277.6).

II.278 Admissible filters and their torsion spectrum

Source 58, delivered Section 3.

Definition II.278.1. An admissible decreasing filter of degree $D \geq 1$ is a real trigonometric polynomial

$$K(t) = \sum_{n=-D}^D c_n e(nt) \quad (\text{II.278.1})$$

such that $0 \leq K(t) \leq 1$, $K(0) = 1$, and

$$c_{-n} = c_n \geq 0, \quad c_0 \geq c_1 \geq \cdots \geq c_D, \quad c_1 > 0. \quad (\text{II.278.2})$$

Coefficients outside $[-D, D]$ are set equal to zero. Define

$$P = \sum_n c_n^m, \quad P_a = \sum_{a|n} c_n^m \quad (a \geq 2), \quad Q = P_2, \quad R_s = \int_{\mathbb{T}} K(t)^s dt. \quad (\text{II.278.3})$$

The normalization implies $\sum_n c_n = 1$. In particular,

$$0 < Q < P < 1, \quad P \leq R_m \leq R_2 \leq 1. \quad (\text{II.278.4})$$

The strict inequality $Q < P$ follows from $c_1 > 0$; $P < 1$ follows because at least three coefficients are positive and $m > 1$. To see $P \leq R_m$, expand the constant coefficient of K^m . Its nonnegative summands include all vectors $\mathbf{n} = t\varepsilon$, whose total weight is P . The last two inequalities follow from $0 \leq K \leq 1$.

Lemma II.278.2 (Order spectrum). *For every integer $a \geq 2$, one has $P_a \leq Q$. More generally, if $a \geq a_0 \geq 2$, then $P_a \leq P_{a_0}$. If $\delta \in H$ has order a , then*

$$\sum_{n \in \mathbb{Z}} c_n^m \mathbf{1}_{n\delta=0} = P_a. \quad (\text{II.278.5})$$

Proof. The exact identity follows from $n\delta = 0$ if and only if $a \mid n$. By symmetry and monotonicity,

$$P_a = c_0^m + 2 \sum_{j \geq 1} c_{aj}^m \leq c_0^m + 2 \sum_{j \geq 1} c_{a_0j}^m = P_{a_0}.$$

All sums are finite because the coefficients are extended by zero. $\square$

This elementary spectrum is the reason small target torsion cannot be discarded. If the order of a discrepancy exceeds D , only frequency zero contributes, and $P_a = c_0^m$. If it has order two, every even frequency contributes. Replacing this contribution by c_0^m would invalidate a uniform theorem for targets containing an element of order two.

Lemma II.278.3 (Positivity ceiling). *For any normalized nonnegative filter with real nonnegative Fourier coefficients, without a monotonicity assumption, and every integer $a \geq 2$,*

$$1 \leq \frac{P}{P_a} \leq a. \quad (\text{II.278.6})$$

Proof. Let $F = K^{*m}$ be the m -fold normalized convolution on $\mathbb{T}$. It is nonnegative, and its Fourier coefficients are c_n^m . The roots of unity identity gives

$$aP_a - P = \sum_{j=1}^{a-1} F(j/a) \geq 0. \quad (\text{II.278.7})$$

Also $P_a \leq P$ because it is a sub-sum of nonnegative terms. $\square$

The order-two case bounds the separation of a single positive filter by two. Increasing the degree cannot make all nonzero discrepancies arbitrarily unlikely relative to zero discrepancy. Independent layers are therefore a natural way to amplify a uniform separation.

II.279 A finite random selector and the good-tuple bound

Source 58, delivered Section 4.

Fix an integer $k \geq 1$ and an integer $M > mD$. For each $\ell = 1, \dots, k$, choose independently a uniform $\chi_\ell \in \widehat{G}$, a uniform $\psi_\ell \in \widehat{H}$, and a shift

$$u_\ell \in \{0, 1/M, \dots, (M-1)/M\}$$

with uniform distribution. Conditional on all these choices, select each *distinct point* $x \in B$ independently with probability

$$p(x) = \prod_{\ell=1}^k K(\chi_\ell(x) + \psi_\ell(\phi(x)) + u_\ell). \quad (\text{II.279.1})$$

The selected set is denoted B' . This is a finite probability space: there are finitely many parameter choices and finitely many subsets, with real weights between zero and one. No limiting random process is needed.

For one layer and an ordered tuple $\mathbf{x}$, expand the *formal occurrence product*. Character orthogonality and the shift average give

$$\begin{aligned} \mathbb{E}_{\chi, \psi, u} \prod_{i=1}^m K(\chi(x_i) + \psi(\phi(x_i)) + u) \\ = \sum_{\mathbf{n} \in [-D, D]^m} \left(\prod_{i=1}^m c_{n_i} \right) \mathbf{1}_{\sum_i n_i = 0} \mathbf{1}_{\sum_i n_i x_i = 0} \mathbf{1}_{\sum_i n_i \phi(x_i) = 0}. \end{aligned} \quad (\text{II.279.2})$$

The first equation here is an equality in $\mathbb{Z}$, while the second and third are in G and H . Indeed, averaging the shifts first gives $\sum_i n_i \equiv 0 \pmod{M}$, and $|\sum_i n_i| \leq mD < M$ upgrades it to the integer equality. In particular, there is no shift-grid aliasing.

Call $\mathbf{n}$ *principal* if $\mathbf{n} = t\varepsilon$ for some integer t . Since $\sum_i \varepsilon_i = 0$, every principal vector passes the shift condition. If $\mathbf{x}$ is respected, it also passes both group conditions, and its total principal weight is P .

Proposition II.279.1 (Good survival, including repetitions). *The selector (II.279.1) satisfies*

$$\mathbb{E} G_m(B', \phi) \geq P^k G_m(B, \phi). \quad (\text{II.279.3})$$

Proof. All the terms in (II.279.2) are nonnegative, so the expected formal product for a respected tuple is at least P in each layer. Independence of the layers gives P^k .

If a point x occurs μ_x times in the tuple, its factor in the formal containment product is $p(x)^{\mu_x}$. Its factor in the actual conditional containment probability is $p(x)$. Since $0 \leq p(x) \leq 1$,

$$\prod_{x \text{ distinct in } \mathbf{x}} p(x) \geq \prod_{i=1}^m p(x_i).$$

Thus the lower bound remains valid with repeated entries. Summing it over the respected tuples proves the proposition. $\square$

The inequality in this proof has a direction. It does not justify an upper bound for repeated bad tuples. Those tuples are treated separately in Section II.281.

II.280 Nonprincipal relations and the domain torsion cost

Source 58, delivered Section 5.

For a finite abelian group G and an integer $d \neq 0$, write

$$G[d] = \{x \in G : dx = 0\}, \quad \tau_G(L) = \max_{1 \leq d \leq L} |G[d]|.$$

If $G = \mathbb{Z}/N\mathbb{Z}$, then $|G[d]| = \gcd(N, d)$, and therefore

$$\tau_{\mathbb{Z}/N\mathbb{Z}}(2D) \leq 2D. \quad (\text{II.280.1})$$

This includes composite N and the trivial group. If $G \cong \bigoplus_{j=1}^s \mathbb{Z}/N_j\mathbb{Z}$, the exact formula is

$$|G[d]| = \prod_{j=1}^s \gcd(N_j, d). \quad (\text{II.280.2})$$

Lemma II.280.1 (Two-relation completion bound). *Let $\mathbf{n} \in [-D, D]^m \cap \mathbb{Z}^m$ be nonprincipal. Then the number of $\mathbf{x} \in B^m$ satisfying*

$$\sum_i \varepsilon_i x_i = 0, \quad \sum_i n_i x_i = 0 \quad (\text{II.280.3})$$

is at most $\tau_G(2D)b^{m-2}$.

Proof. Some two-column minor

$$d = \varepsilon_i n_j - \varepsilon_j n_i$$

is nonzero. Otherwise n_i/ε_i would have the same integer value for every i , making $\mathbf{n}$ principal. Also $0 < |d| \leq 2D$. Fix all entries except x_i, x_j . The first equation in (II.280.3) expresses one of these as an affine function of the other, because its coefficient is ± 1 . Substitution into the second leaves an equation of the form $dx = a$. Every nonempty fibre of multiplication by d is a coset of $G[d]$. Thus there are at most $|G[d]| \leq \tau_G(2D)$ completions for each of the b^{m-2} choices. $\square$

Proposition II.280.2 (Weighted nonprincipal error). *In the k -layer Fourier expansion, the total contribution from arrays containing a nonprincipal layer, summed over any subcollection of additive tuples with distinct entries, is at most*

$$\tau_G(2D)(R_m^k - P^k)b^{m-2}. \quad (\text{II.280.4})$$

Proof. Before imposing either group condition, the total coefficient weight of a single layer satisfying the integer zero-sum condition is

$$\sum_{\sum_i n_i = 0} \prod_i c_{n_i} = R_m.$$

The principal part of this weight is P . Therefore the weight of all k -layer arrays with a nonprincipal layer is exactly $R_m^k - P^k$.

For each such array choose one nonprincipal layer. Every tuple contributing to that array satisfies the two relations in Lemma II.280.1. Ignore all further group equations, including the target equations, and use the completion bound. Sum over the arrays with their nonnegative weights. Distinctness only reduces the number counted. This gives (II.280.4). $\square$

This argument pays neither for the number of coefficient vectors nor for a union bound over the layers. In particular, the error contains $R_m^k - P^k$, rather than a constant independent of k or an extra factor k . This decay is useful when one amplifies a fixed filter many times.

II.280.1 A refined coefficient-by-coefficient torsion profile

For a nonprincipal vector define

$$\tau(\mathbf{n}) = \min_{\varepsilon_i n_j - \varepsilon_j n_i \neq 0} |G[\varepsilon_i n_j - \varepsilon_j n_i]|, \quad W_G(K) = \sum_{\substack{\sum_i n_i = 0 \\ \mathbf{n} \text{ nonprincipal}}} \tau(\mathbf{n}) \prod_i c_{n_i}. \quad (\text{II.280.5})$$

The same proof, classified by the *first* nonprincipal layer, replaces (II.280.4) by

$$W_G(K) \frac{R_m^k - P^k}{R_m - P} b^{m-2}. \quad (\text{II.280.6})$$

Indeed, the first $\ell - 1$ layers have total weight $P^{\ell-1}$, layer ℓ contributes $W_G(K)$, and the remaining layers have weight $R_m^{k-\ell}$. Summing the geometric progression proves the formula. Here $R_m > P$: the positive coefficients c_0, c_1, c_{-1} supply a nonprincipal zero-sum vector for $m \geq 4$.

There is a related exact ambient calculation. The integer matrix with rows ε and $\mathbf{n}$ has Smith invariants $(1, g_{\mathbf{n}})$, where $g_{\mathbf{n}}$ is the greatest common divisor of its nonzero two-column minors. Hence its kernel on G^m has size

$$|G|^{m-2} |G[g_{\mathbf{n}}]|. \quad (\text{II.280.7})$$

To see this, elementary unimodular row and column operations reduce the matrix to its Smith form; the first invariant is one because the first row contains a unit. On G , the resulting two diagonal equations have respectively one and $|G[g_{\mathbf{n}}]|$ choices, with $m - 2$ free coordinates.

The distinction between (II.280.6) and (II.280.7) matters. Smith operations can mix the coordinates and need not preserve B^m . The gcd of all minors therefore cannot simply replace the chosen-minor fibre bound for an arbitrary subset B .

II.281 Repeated bad tuples and their decaying error

Source 58, delivered Section 6.

Lemma II.281.1 (Distinct-point Hölder bound). *If a tuple has $d \geq 2$ distinct entries, its probability of containment in the selector is at most $R_d^k \leq R_2^k$.*

Proof. Fix one layer and its characters. Let the phases of the d distinct points be $\theta_1, \dots, \theta_d$. Hölder's inequality for the finite shift average gives

$$\mathbb{E}_u \prod_{j=1}^d K(u + \theta_j) \leq \prod_{j=1}^d (\mathbb{E}_u K(u + \theta_j)^d)^{1/d}.$$

The translated polynomial $K(u + \theta_j)^d$ has degree at most $dD \leq mD < M$. Its grid average equals its constant coefficient, which is R_d . Thus the right side is R_d . Averaging over the characters preserves the bound. Independence of the layers gives R_d^k , and $R_d \leq R_2$ follows from $0 \leq K \leq 1$ and $d \geq 2$. $\square$

Corollary II.281.2 (Collision error). *The expected number of repeated bad tuples retained in B' is at most*

$$\binom{m}{2} R_2^k b^{m-2}. \quad (\text{II.281.1})$$

Proof. Every bad tuple has at least two distinct entries. Apply Lemma II.281.1 and then the collision count in Lemma II.277.1. $\square$

When the collision profile is known, one can retain more information. Let C_d be the number of bad additive tuples with exactly d distinct entries. The exact improvement is

$$\sum_{d=2}^{m-1} C_d R_d^k. \quad (\text{II.281.2})$$

This may be substantially below (II.281.1). A purely combinatorial bound on C_d is also possible: partition the m positions into d cells, and give a cell the sum of its ε_i coefficients. If all these sums vanish, every tuple with that pattern is respected. Otherwise, solving for a cell with nonzero coefficient a gives at most $|G[a]|b^{d-1}$ assignments. Sum the best such bound over the relevant set partitions. No such refinement is needed for Theorem A.

II.282 The expectation theorem and a simultaneous selection certificate

Source 58, delivered Section 7.

Theorem II.282.1 (Central expectation estimate). *Let $m = 2r \geq 4$, let G, H be finite abelian groups, and let $B \subseteq G$ have cardinality $b > 0$. For an admissible decreasing filter K and $k \geq 1$, the selector of Section II.279 satisfies*

$$\mathbb{E}G_m(B', \phi) \geq P^k G_m(B, \phi), \quad (\text{II.282.1})$$

$$\mathbb{E}U_m(B', \phi) \leq Q^k U_m(B, \phi) + A_k b^{m-2}, \quad (\text{II.282.2})$$

where

$$A_k = \tau_G(2D)(R_m^k - P^k) + \binom{m}{2} R_2^k \leq \left(\tau_G(2D) + \binom{m}{2} \right) R_2^k. \quad (\text{II.282.3})$$

Proof. The first inequality is Proposition II.279.1. For the second, first consider bad tuples with distinct entries, so actual survival equals the formal occurrence product. If the discrepancy has order $a \geq 2$, an all-principal array contributes $P_a^k \leq Q^k$. Summed over such tuples this is at most $Q^k U_m(B, \phi)$. The arrays with a nonprincipal layer contribute at most (II.280.4). Finally add the repeated bad tuples, using Corollary II.281.2. This proves (II.282.2). The simpler upper bound on A_k uses $R_m \leq R_2$ and discards the nonpositive term $-\tau_G(2D)P^k$. $\square$

The expectations must be converted into a single subset having both the desired quality and the desired mass. Separate applications of averaging to two inequalities would not establish this.

Theorem II.282.2 (Exact selection certificate). *Fix $0 < \eta < 1$ and abbreviate $G_m(B, \phi) = g$, $U_m(B, \phi) = u$. Define*

$$\Psi_k = \eta P^k g - (1 - \eta) Q^k u - (1 - \eta) A_k b^{m-2}. \quad (\text{II.282.4})$$

If $\Psi_k > 0$, some subset $B' \subseteq B$ satisfies

$$G_m(B', \phi) \geq (1 - \eta) T_m(B'), \quad T_m(B') \geq \frac{\Psi_k}{\eta} > 0. \quad (\text{II.282.5})$$

Proof. For each outcome of the selector, put

$$F = \mathbf{G}_m(B', \phi) - (1 - \eta)\mathbf{T}_m(B') = \eta\mathbf{G}_m(B', \phi) - (1 - \eta)\mathbf{U}_m(B', \phi).$$

The expectation theorem gives $\mathbb{E}F \geq \Psi_k$. Thus some outcome has $F \geq \Psi_k > 0$, which proves the quality assertion. For that same outcome,

$$F = \eta\mathbf{T}_m(B') - \mathbf{U}_m(B', \phi) \leq \eta\mathbf{T}_m(B'),$$

and the mass assertion follows. $\square$

The exact certificate can use observed counts g, u and the refined errors (II.280.6) or (II.281.2). The following consequence depends only on a lower bound for g .

Theorem II.282.3 (Polynomial restriction from a spectral gap). *Suppose $0 < \alpha, \eta \leq 1$ and $\mathbf{G}_m(B, \phi) \geq \alpha b^{m-1}$. Set*

$$t = \alpha\eta, \quad \mathcal{R} = \frac{P}{Q} > 1, \quad k = \left\lceil \frac{\log(2/t)}{\log \mathcal{R}} \right\rceil. \quad (\text{II.282.6})$$

If

$$b \geq \frac{4A_k}{tP^k}, \quad (\text{II.282.7})$$

then a subset $B' \subseteq B$ satisfies

$$\mathbf{G}_m(B', \phi) \geq (1 - \eta)\mathbf{T}_m(B'), \quad (\text{II.282.8})$$

$$\mathbf{T}_m(B') \geq \frac{\alpha P^k}{4} b^{m-1}. \quad (\text{II.282.9})$$

A simpler sufficient size condition is

$$b \geq 4 \left(\tau_G(2D) + \binom{m}{2} \right) t^{-1} \left(\frac{R_2}{P} \right)^k. \quad (\text{II.282.10})$$

Proof. If $\eta = 1$, take $B' = B$; the required mass follows from the assumption on $\mathbf{G}_m$. Suppose $0 < \eta < 1$. The choice of k gives $Q^k \leq (t/2)P^k$. Since $u \leq \mathbf{T}_m(B) \leq b^{m-1}$,

$$\begin{aligned} \Psi_k &\geq tP^k b^{m-1} - (1 - \eta) \frac{t}{2} P^k b^{m-1} - (1 - \eta) A_k b^{m-2} \\ &\geq \frac{t}{2} P^k b^{m-1} - (1 - \eta) A_k b^{m-2}. \end{aligned}$$

Under (II.282.7), the final error is at most $tP^k b^{m-1}/4$. Thus $\Psi_k \geq tP^k b^{m-1}/4 > 0$. Apply Theorem II.282.2 and divide by η . The coarse condition follows from (II.282.3). $\square$

To exhibit the powers explicitly, define

$$a_K = \frac{-\log P}{\log \mathcal{R}}, \quad z_K = \frac{\log(R_2/P)}{\log \mathcal{R}}, \quad C_G = \tau_G(2D) + \binom{m}{2}. \quad (\text{II.282.11})$$

Since $0 < P < 1$ and $R_2/P \geq 1$, the ceiling gives

$$P^k \geq P(t/2)^{a_K}, \quad (R_2/P)^k \leq (R_2/P)(2/t)^{z_K}. \quad (\text{II.282.12})$$

We obtain the following power-law statement with explicit constants.

Corollary II.282.4. *Under the preceding hypotheses, the sufficient threshold*

$$b \geq 4C_G \frac{R_2}{P} 2^{z_K} (\alpha\eta)^{-(z_K+1)} \quad (\text{II.282.13})$$

guarantees the quality conclusion and

$$\mathsf{T}_m(B') \geq \frac{P}{4 \cdot 2^{a_K}} \alpha^{a_K+1} \eta^{a_K} b^{m-1}. \quad (\text{II.282.14})$$

For fixed m and a fixed filter, these are polynomial dependencies on the two accuracy parameters. In a cyclic domain, every constant is independent of the modulus. For general finite abelian domains the torsion factor is genuine: for example, $\tau_{(\mathbb{Z}/2\mathbb{Z})^s}(2D) = 2^s$ when $D \geq 1$. In such a family the sufficient condition can become vacuous. We do not claim a dimension-independent theorem for all finite abelian domains.

II.283 Powered cosine filters and the certified exponent 54

Source 58, delivered Section 8.

For an integer $q \geq 1$ use

$$K_q(t) = \cos^{2q}(\pi t), \quad c_n = 4^{-q} \binom{2q}{q+n} \quad (|n| \leq q). \quad (\text{II.283.1})$$

This kernel is already used in the prime-field refinement of repository source 04 [236]. It satisfies all the conditions of Definition II.278.1, has degree $D = q$, and has exact moments

$$R_s = 4^{-qs} \binom{2qs}{qs}. \quad (\text{II.283.2})$$

The present argument chooses q as a function of the relation length and then amplifies the same filter; q need not increase as η decreases.

II.283.1 An elementary uniform $O(m \log m)$ bound

Lemma II.283.1. *For every integer $q \geq 1$,*

$$4^{-q} \binom{2q}{q} \geq \frac{1}{2\sqrt{q}}. \quad (\text{II.283.3})$$

Proof. Equality holds at $q = 1$. The ratio of the left sides at $q + 1$ and q is $(2q + 1)/(2q + 2)$, which is at least $\sqrt{q/(q + 1)}$ after squaring. Induction proves the claim. $\square$

Proposition II.283.2 (A uniform spectral gap). *For $m \geq 4$ and the choice $q = m$,*

$$\frac{P}{Q} > \frac{3}{2}, \quad P > \frac{5}{3} (2\sqrt{m})^{-m}. \quad (\text{II.283.4})$$

Consequently

$$a_K \leq \frac{m \log(2\sqrt{m}) - \log(5/3)}{\log(3/2)} = O(m \log m). \quad (\text{II.283.5})$$

Proof. Let $c_0 = 4^{-m} \binom{2m}{m}$. Since

$$\left(\frac{c_1}{c_0}\right)^m = \left(\frac{m}{m+1}\right)^m > e^{-1} > \frac{1}{3},$$

we have $P/c_0^m > 5/3$. Lemma II.283.1 gives the claimed lower bound on P .

For $1 \leq n \leq m$, write

$$\frac{c_n}{c_0} = \prod_{j=1}^n \frac{m+1/2-(j-1/2)}{m+1/2+(j-1/2)}.$$

The inequality $\log((1-u)/(1+u)) \leq -2u$ for $0 < u < 1$ yields

$$\frac{c_n}{c_0} \leq \exp\left(-\frac{n^2}{m+1/2}\right).$$

Thus, extending the sum beyond the support,

$$\frac{Q}{c_0^m} \leq 1 + 2 \sum_{j \geq 1} \exp\left(-\frac{4mj^2}{m+1/2}\right) < 1 + 2 \sum_{j \geq 1} e^{-3j^2} < \frac{10}{9}.$$

For the last inequality, $j^2 \geq 1 + 3(j-1)$ bounds the sum by $e^{-3}/(1-e^{-9})$, and $e^3 > 20$ suffices. Therefore $P/Q > (5/3)/(10/9) = 3/2$. Substitute in (II.282.11). $\square$

Combining this proposition with Corollary II.282.4 proves all-modulus polynomial restriction for every balanced even arity. The bound $z_K \leq a_K$ follows from $R_2 \leq 1$, so the sufficient threshold also has exponent $O(m \log m)$. Theorem B will identify the best possible leading retention exponent within the filter class.

II.283.2 An exact certificate at $m = 16$

Choose $q = 25$, and introduce the integers

$$A = \sum_{n=-25}^{25} \binom{50}{25+n}^{16}, \quad B_{\text{ev}} = \sum_{\substack{-25 \leq n \leq 25 \\ 2|n}} \binom{50}{25+n}^{16}. \quad (\text{II.283.6})$$

Then

$$P = \frac{A}{2^{800}}, \quad Q = \frac{B_{\text{ev}}}{2^{800}}, \quad R_2 = \frac{\binom{100}{50}}{2^{100}}. \quad (\text{II.283.7})$$

The following five inequalities are exact integer certificates. They are verified without floating point in the accompanying program `code/verify_constants.py`; Appendix V.220 also gives the complete essential computation.

Integer inequality	Consequence
$A \geq 2^{750}$	$P \geq 2^{-50}$
$250A \geq 481B_{\text{ev}}$	$\mathcal{R} = P/Q \geq 481/250$
$481^{53} \geq 2^{50} 250^{53}$	$\mathcal{R}^{53} \geq 2^{50}$
$481^{50} \geq 2^{47} 250^{50}$	$\mathcal{R}^{50} \geq 2^{47}$
$\binom{100}{50} \leq 2^{97}$	$R_2 \leq 2^{-3}$

Proof of Theorem A for finite targets. Set $t = \alpha\eta$ and $k = \lceil \log_{\mathcal{R}}(2/t) \rceil$. In a cyclic group $\tau_G(50) \leq 50$ and $\binom{16}{2} = 120$, so (II.282.10) is implied by

$$b \geq 680t^{-1}(R_2/P)^k. \quad (\text{II.283.8})$$

The certificates give $R_2/P \leq 2^{47}$ and $\log_{\mathcal{R}}(2^{47}) \leq 50$. Hence

$$(R_2/P)^k \leq 2^{47}(2/t)^{50},$$

so the right side of (II.283.8) is at most $680 \cdot 2^{97}t^{-51}$. Since $680 < 2^{11}$ and $0 < t \leq 1$,

$$680 \cdot 2^{97}t^{-51} \leq 2^{108}t^{-54} = (4/t)^{54}. \quad (\text{II.283.9})$$

Thus the hypothesis of Theorem A ensures the size condition of Theorem II.282.3.

For the retention, the other certificates imply

$$a_K \leq 53, \quad P^k \geq P(t/2)^{a_K} \geq 2^{-50}(t/2)^{53}.$$

Consequently

$$\frac{\alpha P^k}{4} \geq 2^{-105} \alpha^{54} \eta^{53} \geq 2^{-108} \alpha^{54} \eta^{54} = \left(\frac{\alpha\eta}{4} \right)^{54}. \quad (\text{II.283.10})$$

Theorem II.282.3 proves the tuple-count conclusions. Finally $\mathsf{T}_{16}(B') \leq |B'|^{15}$ gives the point-count corollary. $\square$

The exact rational values, rather than their rounded certificates, give

$$\begin{aligned} P &\approx 1.4263285575 \times 10^{-15}, & \mathcal{R} &\approx 1.9247171505, \\ a_K &\approx 52.2064261968, & z_K &\approx 48.3411887375. \end{aligned} \quad (\text{II.283.11})$$

These decimals are descriptive, not the basis of the integer theorem. For a precise nonintegral-exponent theorem, use the exact definitions (II.283.7) in (II.282.13) and (II.282.14). Degree 25 was selected by a finite numerical search within the powered cosine family. We do not claim that it optimizes over all admissible filters, or even over all degrees without a certified global search.

II.284 Further generalizations of the restriction estimate

Source 58, delivered Section 9.

II.284.1 Targets of arbitrary cardinality

The finite-target hypothesis can be removed for the restriction conclusions, because only finitely many target discrepancies are relevant.

Lemma II.284.1 (Finite quotient preserving all relevant relations). *Let B be finite and $\phi : B \rightarrow H$, where H is any abelian group. For a fixed balanced arity $m = 2r$, there are a finite abelian group H_0 and a map $\phi_0 : B \rightarrow H_0$ such that every ordered tuple in B^m has zero balanced target discrepancy under ϕ if and only if it does under ϕ_0 .*

Proof. The subgroup generated by $\phi(B)$ is finitely generated, hence is isomorphic to $\mathbb{Z}^s \oplus F$ with F finite. Consider the finite set of all balanced m -term discrepancies of points of B . Choose an integer $L \geq 2$ larger than the absolute value of every coordinate of every free part in this set. Reduction of $\mathbb{Z}^s$ modulo L , keeping F unchanged, kills none of these finitely many nonzero discrepancies: an element with a nonzero free coordinate has a nonzero coordinate modulo L , and an element with zero free part has its nonzero torsion component preserved. Compose ϕ with this finite quotient. Zero discrepancies remain zero and all the relevant nonzero discrepancies remain nonzero. $\square$

Corollary II.284.2. *Theorem A and the restriction conclusions of Theorem II.282.3 and Corollary II.282.4 hold for maps to arbitrary abelian target groups, with unchanged constants.*

Proof. Apply Lemma II.284.1 and the finite-target theorem. The same quotient preserves the relevant tuple counts on every subset of B , so the selected subset works for the original map. $\square$

II.284.2 Using the actual discrepancy-order distribution

For a finite target let u_a be the number of bad tuples in B^m whose discrepancy has order a . The proof of Theorem II.282.1 gives the stronger estimate

$$\mathbb{E}U_m(B', \phi) \leq \sum_{a \geq 2} P_a^k u_a + A_k b^{m-2}. \quad (\text{II.284.1})$$

It is harmless that the sum includes repeated tuples, since those only enlarge the principal upper bound and their actual survival is already charged to the collision error. Only finitely many u_a are nonzero.

If every nonzero discrepancy has order at least a_0 , monotonicity gives $P_a \leq P_{a_0}$, and all the restriction theorems hold with $Q = P_{a_0}$. In particular, a finite target of odd order uses $a_0 \geq 3$. More generally, for a nontrivial finite target its smallest prime divisor is an available choice for a_0 . If all nonzero discrepancies have order greater than D , one can use $Q = c_0^m$. This explains why a large-characteristic argument can have a stronger single-filter separation than an argument uniform over targets containing elements of order two.

Coefficient monotonicity can also be relaxed. For a nonnegative coefficient filter with $c_1 > 0$, use $Q_* := \max_{a \geq 2} P_a$ instead. This is a maximum over finitely many distinct possibilities, since $P_a = c_0^m$ for $a > D$, and $Q_* < P$ because every P_a omits $c_1^m + c_{-1}^m$. The same proof applies. The decreasing condition conveniently identifies the worst order as two.

II.284.3 Different filters in different layers

Let layer ℓ use an admissible filter K_ℓ of degree D_ℓ , with parameters $P_\ell, Q_\ell, R_{s,\ell}$. Take $M > m \max_\ell D_\ell$ and set $D_* := \max_\ell D_\ell$. Then

$$\begin{aligned} \mathbb{E}G_m(B', \phi) &\geq \left(\prod_\ell P_\ell \right) G_m(B, \phi), \\ \mathbb{E}U_m(B', \phi) &\leq \left(\prod_\ell Q_\ell \right) U_m(B, \phi) \\ &\quad + \left\{ \tau_G(2D_*) \left(\prod_\ell R_{m,\ell} - \prod_\ell P_\ell \right) + \binom{m}{2} \prod_\ell R_{2,\ell} \right\} b^{m-2}. \end{aligned} \quad (\text{II.284.2})$$

The proof is identical: nonprincipal arrays have the difference of the two displayed product weights, while the distinct-point Hölder bound multiplies across layers. This flexibility may improve a finite-size certificate even when it does not improve the asymptotic principal exponent, as the next section shows.

II.285 Optimal leading growth of the Fourier-filter exponent

Source 58, delivered Section 10.

In this section m is any integer at least two, not necessarily even. Let $\mathcal{K}$ be the class of finite trigonometric polynomials

$$K(t) = \sum_n c_n e(nt), \quad 0 \leq K \leq 1, \quad K(0) = 1, \quad c_{-n} = c_n \geq 0.$$

For a fixed integer $a \geq 2$, use the notation of Theorem B and let

$$I_{m,a} = \inf_{K \in \mathcal{K}} \mathcal{A}_{m,a}(K). \quad (\text{II.285.1})$$

We prove both sides of the asymptotic, uniformly over all finite degrees in the lower bound.

II.285.1 A coefficient bound from the mean of the filter

Put $s = c_0 = \int_{\mathbb{T}} K$. Continuity and $K(0) = 1$ give $s > 0$. If $s = 1$, then $K \equiv 1$, so the exponent is infinite by convention. We may assume $0 < s < 1$. The integral formula gives $c_n \leq s$, and $\sum_n c_n = 1$. Consequently

$$P_m \leq s^{m-1} < 1, \quad Q_{m,a} \geq s^m > 0. \quad (\text{II.285.2})$$

Lemma II.285.1 (Rearrangement bound for a nonzero coefficient). *For every nonzero integer n ,*

$$c_n \leq \frac{\sin(\pi s)}{\pi} = s\beta(s), \quad \beta(s) := \frac{\sin(\pi s)}{\pi s}. \quad (\text{II.285.3})$$

Moreover

$$0 < \beta(s) \leq 1 - \frac{2s^2}{3} \leq \exp(-2s^2/3) \quad (0 < s < 1). \quad (\text{II.285.4})$$

Proof. Let $h(t) = \cos(2\pi nt)$. Its distribution on $\mathbb{T}$ is the same as that of $\cos(2\pi t)$. Let E be its top level set of measure s , so its boundary level is $\cos(\pi s)$. Pointwise,

$$(\mathbf{1}_E - K)(h - \cos(\pi s)) \geq 0.$$

Since $\int(\mathbf{1}_E - K) = 0$, integration gives

$$c_n = \int K(t)h(t) dt \leq \int_E h(t) dt = \int_{-s/2}^{s/2} \cos(2\pi t) dt = \frac{\sin(\pi s)}{\pi}.$$

This proves the coefficient bound directly, without an extremal regularity assumption on K .

For $0 \leq v \leq 1$, concavity of sine on $[0, \pi/2]$ gives $\sin(\pi v/2) \geq v$, hence $\cos(\pi v) \leq 1 - 2v^2$. Integrating with $v = su$ for $0 \leq u \leq 1$ yields

$$\beta(s) = \int_0^1 \cos(\pi su) du \leq 1 - 2s^2/3.$$

The exponential bound follows from $1 - x \leq e^{-x}$. □

II.285.2 The universal lower bound

If $P_m = Q_{m,a}$, the lower bound is automatic. Otherwise let $\mathcal{R}_a = P_m/Q_{m,a} > 1$. By Lemma II.278.3, $\mathcal{R}_a \leq a$. Equation (II.285.2) gives

$$\mathcal{A}_{m,a}(K) \geq \frac{-(m-1)\log s}{\log a}. \quad (\text{II.285.5})$$

On the other hand, Lemma II.285.1 implies

$$\begin{aligned} \mathcal{R}_a - 1 &= \frac{P_m - Q_{m,a}}{Q_{m,a}} \\ &\leq \frac{\sum_{n \neq 0} c_n^m}{s^m} \leq \frac{1-s}{s} \beta(s)^{m-1} \leq \frac{1-s}{s} \exp\left(-\frac{2(m-1)s^2}{3}\right). \end{aligned} \quad (\text{II.285.6})$$

Use $\log \mathcal{R}_a \leq \mathcal{R}_a - 1$ and $-\log s \geq 1 - s$ to get a complementary bound,

$$\mathcal{A}_{m,a}(K) \geq (m-1)s \exp\left(\frac{2(m-1)s^2}{3}\right). \quad (\text{II.285.7})$$

For sufficiently large m , put $s_m = \sqrt{3 \log m / m} < 1$. If $s \leq s_m$, equation (II.285.5) yields

$$\mathcal{A}_{m,a}(K) \geq \frac{m-1}{2 \log a} (\log m - \log \log m - \log 3). \quad (\text{II.285.8})$$

If $s > s_m$, the right side of (II.285.7) is increasing in s , and therefore

$$\mathcal{A}_{m,a}(K) \geq (m-1) \sqrt{\frac{3 \log m}{m}} m^{2-2/m}. \quad (\text{II.285.9})$$

For fixed a , this is eventually larger than (II.285.8). Thus, uniformly over $K \in \mathcal{K}$,

$$I_{m,a} \geq \frac{m \log m}{2 \log a} - O_a(m \log \log m). \quad (\text{II.285.10})$$

The two branches express an obstruction: a narrow filter has a small mean and loses many good tuples, whereas a filter with larger mean has little nonzero-frequency mass capable of separating a nonzero discrepancy from zero.

II.285.3 An explicit family attaining the leading constant

Take

$$q = \left\lceil m(\log m)^2 \right\rceil, \quad K(t) = \cos^{2q}(\pi t), \quad J = \left\lfloor \sqrt{q/m} \right\rfloor. \quad (\text{II.285.11})$$

For sufficiently large m , $J \leq q/4$. For $1 \leq n \leq J$,

$$\frac{c_n}{c_0} = \prod_{j=1}^n \left(1 - \frac{2j-1}{q+j}\right).$$

The fractions lie in $[0, 1/2]$, so $\log(1-x) \geq -2x$ gives

$$\log \frac{c_n}{c_0} \geq -\frac{4}{q} \sum_{j=1}^n j \geq -\frac{4n^2}{q}. \quad (\text{II.285.12})$$

Hence $(c_n/c_0)^m \geq e^{-4}$ when $|n| \leq J$, and

$$\frac{P_m}{c_0^m} \geq 1 + 2Je^{-4}. \quad (\text{II.285.13})$$

Let $w_n = c_n^m$. It is even and nonincreasing for $n \geq 0$. Group the positive indices into blocks of length a . Each block ending at aj has every value at least w_{aj} , so

$$\sum_{n \geq 1} w_n \geq a \sum_{j \geq 1} w_{aj}.$$

Together with (II.278.7), this proves

$$0 \leq aQ_{m,a} - P_m \leq (a-1)c_0^m. \quad (\text{II.285.14})$$

Write $d = (aQ_{m,a} - P_m)/P_m$. By (II.285.13),

$$0 \leq d \leq \frac{a-1}{1+2Je^{-4}} = O_a(1/\log m).$$

It follows that

$$\log(P_m/Q_{m,a}) = \log a - \log(1+d) = \log a - O_a(1/\log m). \quad (\text{II.285.15})$$

Finally, Lemma II.283.1 and $P_m \geq c_0^m$ give

$$-\log P_m \leq m \log(2\sqrt{q}) = \frac{m}{2} \log m + m \log \log m + O(m). \quad (\text{II.285.16})$$

Dividing (II.285.16) by (II.285.15) proves

$$I_{m,a} \leq \frac{m \log m}{2 \log a} + O_a(m \log \log m). \quad (\text{II.285.17})$$

Together with (II.285.10), this proves Theorem B, including the asserted additive error term.

II.285.4 What the optimality statement covers

The theorem optimizes the loss in the principal good-to-bad separation. For identical layers this loss is exactly $\mathcal{A}_{m,a}(K)$: achieving separation L costs a principal good-survival factor approximately $L^{-\mathcal{A}_{m,a}(K)}$. It does not optimize the finite-size error or the ceiling in the number of layers. Nor is it a lower bound on the actual survival of every input set: particular inputs can benefit from additional Fourier terms or structure.

Allowing a fixed list of different independent filters cannot improve the optimal principal exponent. If all the denominators are positive, then

$$\frac{\sum_{\ell} -\log P_{m,\ell}}{\sum_{\ell} \log(P_{m,\ell}/Q_{m,a,\ell})} \geq I_{m,a}, \quad (\text{II.285.18})$$

because the left side is a weighted average of the individual exponents, with nonnegative denominator weights. Layers with ratio one contribute no separation and cannot improve this inequality. Adaptive choices depending on earlier outcomes, filters without nonnegative Fourier coefficients, and structural extraction arguments are not covered by this lower bound.

For $a = 2$, the theorem explains a difference from a large-characteristic prime-field argument. The latter can exclude all nonzero principal frequencies for every bad discrepancy in a chosen degree range, whereas a uniform order-two argument has separation at most two per layer. There is consequently no conflict between the $m \log m$ barrier here and the linear-in- m prime-field exponents already present in sources 04 and 17.

II.286 Moment lifting and the Gowers–ProveIt interfaces

Source 58, delivered Section 11.

We return to balanced even arity $m = 2r \geq 4$. The restriction theorem accepts a density of respected m -tuples. In applications the input is often a density of respected quadruples. The following standard interpolation, already used in repository sources 04 and 17, performs the conversion without an extra ambient-density loss.

Proposition II.286.1 (Graph moment lifting). *For a nonempty finite set A in a finite abelian group V , let $E_r(A)$ count ordered solutions to $a_1 + \cdots + a_r = a_{r+1} + \cdots + a_{2r}$ in A^{2r} . Then*

$$E_r(A) \geq \frac{E_2(A)^{r-1}}{|A|^{r-2}} \quad (r \geq 2). \quad (\text{II.286.1})$$

In particular, if $G_4(B, \phi) \geq \gamma b^3$ with $0 < \gamma \leq 1$, then

$$G_{2r}(B, \phi) \geq \gamma^{r-1} b^{2r-1}. \quad (\text{II.286.2})$$

Proof. Orthogonality gives

$$E_j(A) = \frac{1}{|V|} \sum_{\chi \in \widehat{V}} |\widehat{\mathbf{1}_A}(\chi)|^{2j}.$$

Apply Hölder to the normalized Fourier average, writing $|\widehat{\mathbf{1}_A}|^4$ as the product of $|\widehat{\mathbf{1}_A}|^{2r/(r-1)}$ and $|\widehat{\mathbf{1}_A}|^{2(r-2)/(r-1)}$. For $r > 2$ this gives

$$E_2(A) \leq E_r(A)^{1/(r-1)} E_1(A)^{(r-2)/(r-1)}.$$

Since $E_1(A) = |A|$, rearrangement proves (II.286.1); $r = 2$ is equality. Apply it to the graph $A = \{(x, \phi(x)) : x \in B\} \subseteq G \times H$, which has size b and $E_j(A) = G_{2j}(B, \phi)$. For arbitrary abelian targets, apply a finite quotient preserving both the quadruple and $2r$ -tuple discrepancies; the same calculation then proves the assertion there as well. $\square$

Corollary II.286.2 (All-modulus quadruple-input restriction). *Let $B \subseteq \mathbb{Z}/N\mathbb{Z}$ and let $\phi : B \rightarrow H$ for any abelian group H . Suppose $0 < \gamma, \eta \leq 1$ and*

$$G_4(B, \phi) \geq \gamma b^3, \quad b \geq \left(\frac{4}{\gamma^7 \eta} \right)^{54}. \quad (\text{II.286.3})$$

Then a subset $B' \subseteq B$ satisfies

$$G_{16}(B', \phi) \geq (1 - \eta) T_{16}(B'), \quad T_{16}(B') \geq \left(\frac{\gamma^7 \eta}{4} \right)^{54} b^{15}. \quad (\text{II.286.4})$$

Proof. Use $r = 8$ in Proposition II.286.1, and then apply Theorem A and Corollary II.284.2 with $\alpha = \gamma^7$. $\square$

The improvement in moment normalization is credited to the prior project work. The new point in this composition is the all-modulus restriction and its explicit constants.

II.286.1 An explicit replacement at Lemma 9.3

Write $b = \beta N$ as in the original lemma. Under the nonvacuous hypotheses one has $0 < \alpha, \beta \leq 1$. A sufficient starting modulus for Theorem A is

$$N_0 = \left\lceil \beta^{-1} \left(\frac{4}{\alpha\eta} \right)^{54} \right\rceil. \quad (\text{II.286.5})$$

For every $N \geq N_0$, the conclusion becomes

$$\mathsf{T}_{16}(B') \geq \left(\frac{\alpha\eta}{4} \right)^{54} \beta^{15} N^{15}, \quad \mathsf{G}_{16}(B', \phi) \geq (1 - \eta) \mathsf{T}_{16}(B'). \quad (\text{II.286.6})$$

Since $0 < \alpha\eta/4 < 1$ and $54 < 2^{19}$, this also implies the retention required by the original and formal Lemma 9.3. This is a mathematical implication of statements; it is not a claim that the new implication has been checked in Lean.

For Corollary 9.4, the quadruple input is $\mathsf{G}_4(B, \phi) \geq \gamma\beta^3 N^3 = \gamma b^3$. We may take

$$N_0 = \left\lceil \beta^{-1} \left(\frac{4}{\gamma^7\eta} \right)^{54} \right\rceil, \quad (\text{II.286.7})$$

with retention $(\gamma^7\eta/4)^{54}\beta^{15}N^{15}$. The printed coefficient contains the additional factor β^6 inside $(\gamma^7\beta^6\eta/4)^{2^{19}}$. The removal of that density factor in the moment input was already established in the prior prime-field refinements; the present corollary makes the resulting restriction available for every modulus.

II.286.2 Formalization and integration plan

At the inspected commit the repository's definitions match the present counts exactly:

Mathematical quantity	Existing Lean interface
$\mathsf{T}_{16}(B)$	<code>additiveTupleCount 8 B</code>
$\mathsf{G}_{16}(B, \phi)$	<code>phiAdditiveTupleCount 8 B phi</code>
$\mathsf{G}_{16} \geq (1 - \eta) \mathsf{T}_{16}$	<code>GammaHomOfOrder 8 B phi (1 - eta)</code>
Lemma 9.3 / Corollary 9.4 statements	<code>Sections08_09.lean</code>
Existing exact companions	<code>lemma_9_3_holds, corollary_9_4_holds in Proofs09Restriction.lean</code>

The formal statements quantify over all nonzero moduli, not just primes. Their premises quantify $\alpha, \beta > 0$ without explicit upper bounds. Those upper bounds follow in a nonempty instance from $b \leq N$ and $\mathsf{G}_{16} \leq b^{15}$. Thus this difference in presentation causes no interface gap. For Corollary 9.4 the definition `GammaAdditive B phi (gamma * beta^3)` is the ambient normalization of the quadruple condition used above.

[write] **7 October 2026; sources 54–59.** *Checked at HEAD fe7165a3c.* `additiveTupleCount`, `phiAdditiveTupleCount`, `GammaHomOfOrder` and `GammaAdditive` are at lines 322, 327, 334 and 306 of `Definitions.lean`; `lemma_9_3` and `corollary_9_4` at lines 77 and 91 of `Sections08_09.lean`; `lemma_9_3_holds` and `corollary_9_4_holds` at lines 376 and 416 of `Proofs09Restriction.lean`. No new statement with exponent 54 has been added to the development.

A useful implementation order is:

1. Prove the two-relation completion lemma for integer coefficients on finite abelian groups, with the cyclic gcd corollary.
2. Formalize the finite shift-grid identity and the distinction between products over positions and products over distinct points.
3. Establish the exact nonprincipal coefficient mass and the repeated-tuple Hölder bound.
4. Prove the signed expectation certificate and instantiate its finite probability space with the cosine coefficients.
5. Certify the five integer inequalities by exact arithmetic, then expose new strengthened statements with exponent 54 and the explicit N_0 ; derive the existing weaker interfaces from them.

No Lean source is included in this package, because a theorem statement with an unproved placeholder would not be a formalization. The accompanying source audit supplies the pinned provenance needed for integration. The existing report can incorporate these sections as a new contribution while preserving the earlier source identifiers and their priority credits.

II.287 Source 57's chapter: Theorem 13.12 in writing, and formally [write]

Added in the tenth write, 7 October 2026.

The next three sections are source 57's Sections 2–4. *Logically they continue the ninth write's Section 13 chapter* (Section II.195) and use source 40's Fejér selection (Section II.164); they are printed here because no number printed earlier may move. Source 57's introduction is Section F.62, its endpoint section is in Section II.236.

What is proved. For $0 < \alpha < 1$, every prime $p \geq N_0(\alpha)$ and $|f| \leq 1$ with $\|f\|_{U^4}^{16} > \alpha$ there are proper progressions P, Q with the same nonzero step and integer length $m \geq p^{2-\alpha-2^{70}}$ (the printed length), $B \subseteq P \times Q$ with $|B| \geq (\alpha/2)^{2^{41}} m^2$ (the printed density exponent is 2^{76}) and one formula $c_{00} + c_{10}h + c_{01}k + c_{11}hk$ at which $|\widehat{\Delta_h \Delta_k f}| \geq \alpha/2$ on B (Theorem F.62.1). The route: Gowers's Proposition 6.1, Corollary 7.6 and Lemma 12.4, imported with their constants, give a separately Freiman frequency graph of density $a^{14,409,429}$, $a = \alpha/2$ (Proposition II.290.1); source 40's one-stage Fejér selection, reproved for the eight-arrangements (Theorem II.289.1), purifies it to density $a^{2,996,201,429}$ with respected proportion $1 - 2^{-44}$ (Proposition II.290.2); the repository's uniform square extraction, imported at its quantifier order (Input theorem II.288.5), gives the length exponent 2^{53} in $(2/\alpha)^\cdot$ (Theorem II.290.3); and the near-maximal theorem (Theorem II.247.5) covers $1 - \alpha \leq 2^{-10}$, where the two exponents need not compare. Source 57 answers the obligation “resolve the printed Theorem 13.12 length exponent” of `FORMALIZATION_STATUS.txt` at its pin in writing.

Relation to the formal statements (checked at HEAD `fe7165a3c`). `theorem_13_12` (Section `ns12_13.lean`, line 718) has been an exact companion since `d48e8af00`, `theorem_13_12_holds` (`Proofs13EndpointPrintedTheorem.lean`, line 73), proved along this route with the repository's own constants (note after Theorem II.290.3): density $(\alpha/2)^{2^{42}}$, not source 57's 2^{41} , which is not formalized; length exponent 2^{53} , and 2^{61} in place of the printed 2^{70} by the repository's own strengthening. The threshold N_{sq} that source 57 leaves inexplicit is made explicit both by source 53 and formally (note after Remark II.290.4).

II.288 Conventions and the exact imported results

Source 57, delivered Section 2.

Throughout the Fourier part, p is prime, $\mathbb{F}_p$ has normalized counting measure, and $e_p(t) = \exp(2\pi it/p)$. We use

$$\widehat{f}(r) = \mathbb{E}_{x \in \mathbb{F}_p} f(x) e_p(-rx), \quad \Delta_h f(x) = f(x+h) \overline{f(x)}.$$

Thus the Fourier transform used in Gowers's paper is $p\widehat{f}$. The source uses the backward derivative $\nabla_h f(x) = f(x) \overline{f(x-h)}$. The identity

$$\nabla_{h_1} \cdots \nabla_{h_s} f(x) = \Delta_{h_1} \cdots \Delta_{h_s} f(x - h_1 - \cdots - h_s)$$

shows that its Fourier coefficient at r differs from ours by the unit factor $e_p(-r(h_1 + \cdots + h_s))$. In particular, the moduli and frequency labels in every imported hypothesis and conclusion agree. For a finite abelian group G we instead index the normalized Fourier transform by its character group $\widehat{G}$. Write $Q_d(f) = \|f\|_{U^d(G)}^{2^d}$. The basic identities needed below are

$$Q_2(f) = \sum_{\chi \in \widehat{G}} |\widehat{f}(\chi)|^4, \quad \sum_{\chi \in \widehat{G}} |\widehat{f}(\chi)|^2 = \mathbb{E}_x |f(x)|^2, \quad (\text{II.288.1})$$

$$Q_{s+2}(f) = \mathbb{E}_{h_1, \dots, h_s} Q_2(\Delta_{h_1} \cdots \Delta_{h_s} f). \quad (\text{II.288.2})$$

All L^r norms are normalized. In particular, $|f| \leq 1$ implies $0 \leq Q_d(f) \leq 1$. The failure of cubic α -uniformity in the pinned catalogue is exactly $Q_4(f) > \alpha$.

An arithmetic progression in $\mathbb{F}_p$ is represented by $P = \{x + jd : 0 \leq j < m\}$, together with its integer length m . It is *proper* when these m elements are distinct. We always require $d \neq 0$. The full field, represented with $m = p$ and $d = 1$, is a proper progression. A bilinear formula in the terminology of Section 13 means

$$\mu(h, k) = c_{00} + c_{10}h + c_{01}k + c_{11}hk. \quad (\text{II.288.3})$$

It need not be additive in either variable. We call this a *multiaffine formula* when distinguishing it from the homogeneous bilinear map chk .

II.288.1 Arrangements and Freiman restrictions

A map $\phi : A \rightarrow \mathbb{F}_p$, where $A \subseteq \mathbb{F}_p^2$, is separately Freiman of order eight if each horizontal and vertical restriction preserves every equality between two sums of eight elements of that fibre, with repetitions permitted.

Definition II.288.1 (Ordered eight-arrangements). Put $\sigma_j = 1$ for $1 \leq j \leq 8$ and $\sigma_j = -1$ for $9 \leq j \leq 16$. An ordered eight-arrangement is a tuple

$$(x_1, \dots, x_{16}, y_1, \dots, y_{16}, h), \quad \sum_{j=1}^{16} \sigma_j x_j = 0.$$

Its 32 labelled vertices are $v_{j,0} = (x_j, y_j)$ and $v_{j,1} = (x_j, y_j + h)$. It lies in A if all these vertices lie in A , and it is *respected* by ϕ if

$$\sum_{j=1}^{16} \sigma_j (\phi(x_j, y_j + h) - \phi(x_j, y_j)) = 0.$$

Denote the total and respected counts by $T(A)$ and $R(A, \phi)$. Degenerate arrangements and repeated physical vertices are counted.

Lemma II.288.2 (A uniform upper count). *If $|A| = \beta p^2$, then $T(A) \leq \beta^{15} p^{32}$.*

Proof. For fixed h, x , let $n_h(x)$ count the y for which both $(x, y), (x, y + h)$ lie in A . Then $n_h(x) \leq p$ and $\sum_x n_h(x) \leq |A|$. For each h , the first fifteen x coordinates determine x_{16} , so the number of arrangements at height h is at most $p(\sum_x n_h(x))^{15} \leq p|A|^{15}$. Summing over h gives $p^2|A|^{15} = \beta^{15} p^{32}$. $\square$

The following two inputs are from Gowers [1], with the normalization and the exact constants recorded here.

Input theorem II.288.3 (One-fibre restriction). *Suppose $g : \mathbb{F}_p \rightarrow \mathbb{C}$ satisfies $|g| \leq 1$, $U \subseteq \mathbb{F}_p$ has density $u > 0$, and $\psi : U \rightarrow \mathbb{F}_p$ obeys $|\widehat{\Delta_k g}(\psi(k))| \geq a$ for every $k \in U$, where $0 < a \leq 1$. There is $V \subseteq U$ of size at least*

$$2^{-1882} a^{9312} u^{1165} p$$

on which ψ is a Freiman homomorphism of order eight.

Derivation from the cited inputs. The sum of the squared *unnormalized* Fourier coefficients is at least $ua^2 p^3$. Proposition 6.1 gives at least $u^4 a^8 p^3$ ψ -additive quadruples. Corollary 7.6 says that a graph on a density- u set with at least $\gamma(up)^3$ additive quadruples has an order-eight restriction of size $2^{-1882} \gamma^{1164} up$. Taking $\gamma = ua^8$ gives the display. These are the original constants; no strengthened Freiman theorem is being assumed. $\square$

Input theorem II.288.4 (Gowers's Lemma 12.4). *Let $|f| \leq 1$, let $A \subseteq \mathbb{F}_p^2$ have density $\beta > 0$, and suppose $|\widehat{\Delta_k \Delta_l f}(\phi(k, l))| \geq a > 0$ on A . Then*

$$R(A, \phi) \geq \beta^{112} a^{336} p^{32}.$$

II.288.2 The uniform common-square interface

The geometric extraction is the substantial inherited part of Section 13. It must be imported with its *uniform density* quantifiers: a lower density bound is fixed before the modulus and before the actual input set. The statement below is precisely the interface supplied by the pinned repository, not an assumption of Theorem 13.12 itself.

Input theorem II.288.5 (Uniform common-square extraction). *For every $0 < \delta \leq 1$ there is an integer $N_{\text{sq}}(\delta)$ such that the following holds for every prime $p \geq N_{\text{sq}}(\delta)$. Suppose $A \subseteq \mathbb{F}_p^2$ has actual density $\beta \geq \delta$, $\phi : A \rightarrow \mathbb{F}_p$ is separately Freiman of order eight, and*

$$R(A, \phi) \geq (1 - 2^{-44})T(A).$$

There are proper progressions P, Q with the same nonzero step and the same integer length m , a set $B \subseteq A \cap (P \times Q)$, and a global formula (II.288.3) agreeing with ϕ on B , such that

$$m \geq p^{\gamma(\delta)}, \quad |B| \geq 2^{-137} \delta^{704} m^2, \quad (\text{II.288.4})$$

$$\gamma(\delta) = \frac{2^{-386} \delta^{1856}}{2^{13} \mathcal{Q}(\delta)}, \quad \mathcal{Q}(\delta) = 2^{220} \delta^{-221}. \quad (\text{II.288.5})$$

The exact Lean declaration is

`section13_complete_square_extraction_uniform_density`

in `Proofs13UniformSquareExtraction.lean` [232]. The exponent formula is in `Proofs13CompleteSquareExtraction.lean`. The conclusion called `BilinearOn` there means agreement with one global formula of the displayed kind. Extending that formula to $P \times Q$ is therefore legitimate, even though agreement with the selected frequencies is needed only on B . Partial maps can be extended arbitrarily off their specified domains when constructing the total map in the Lean context; every relevant arrangement has all vertices inside that domain.

Lemma II.288.6 (Absorbing the prefactor). *For $0 < \delta \leq 1$, $\gamma(\delta) \geq 2^{-15Q(\delta)}$.*

Proof. Put $x = \delta^{-1} \geq 1$. The elementary estimate $\log_2 x \leq 2x$ follows from $\log x \leq x$ and $\log 2 > 1/2$. Since $Q(\delta) \geq 3712\delta^{-1}$ and $Q(\delta) \geq 386$, we obtain

$$386 + 1856 \log_2(\delta^{-1}) \leq 2Q(\delta).$$

Taking binary logarithms in (II.288.5) proves the claim. The same estimate is already proved in `Proofs13ExplicitFourierExponent.lean`; the short calculation is included to expose the quantitative dependence. $\square$

II.289 Fejér purification for the Section 13 arrangements

Source 57, delivered Section 3.

We reprove the precise specialization of the earlier Fejér selection theorem [233] that will be used. This selection theorem is an inherited ingredient. Its propagation through the common-square interface, and the endpoint argument that completes the parameter range, are the refinements developed here.

Theorem II.289.1 (Specialized Fejér purification). *Let $0 < \beta, \rho, \eta \leq 1$, let $A \subseteq \mathbb{F}_p^2$ have density β , and let $\phi : A \rightarrow \mathbb{F}_p$ satisfy $R(A, \phi) \geq \rho\beta^{15}p^{32}$. If*

$$p \geq 2\beta^{-15} \left(\frac{18 \cdot 33^2}{(\rho\eta)^2} \right)^{32}, \quad (\text{II.289.1})$$

there is $A' \subseteq A$ with

$$R(A', \phi) \geq \frac{\rho^{32}\eta^{31}}{3^{31}33^{32}} \beta^{15}p^{32}, \quad (\text{II.289.2})$$

$$R(A', \phi) \geq \frac{1}{1+\eta} T(A') \geq (1-\eta)T(A'). \quad (\text{II.289.3})$$

II.289.1 The relation isolated by one scalar feature

Put $\lambda_{j,0} = -\sigma_j$, $\lambda_{j,1} = \sigma_j$, and use the single feature $F(x, y) = xy$. There are 32 free arrangement parameters: eliminate x_{16} using $x_{16} = \sum_{j < 16} \sigma_j x_j$. For a vector of coefficients $u = (u_{j,e}) \in \mathbb{F}_p^{32}$, the feature relation is the polynomial

$$\mathcal{P}_u = \sum_{j=1}^{16} x_j((u_{j,0} + u_{j,1})y_j + u_{j,1}h) \quad (\text{II.289.4})$$

after that elimination.

Lemma II.289.2 (Exactly one universal relation). *The polynomial $\mathcal{P}_u$ is identically zero if and only if $u = c\lambda$ for some $c \in \mathbb{F}_p$. Every other $\mathcal{P}_u$ is a nonzero polynomial of total degree at most two.*

Proof. For $j < 16$, the coefficients involving y_j force $u_{j,0} + u_{j,1} = 0$. The coefficient of y_{16} is $(u_{16,0} + u_{16,1}) \sum_{j < 16} \sigma_j x_j$, so its vanishing forces the same identity for $j = 16$. What remains is

$$h \sum_{j < 16} (u_{j,1} + \sigma_j u_{16,1}) x_j.$$

Thus $u_{j,1} = \sigma_j c$ with $c = -u_{16,1}$, and the other coordinates equal $-\sigma_j c$. Conversely these coefficients give $ch \sum_j \sigma_j x_j = 0$. $\square$

Let $L = \lceil 66/(\rho\eta) \rceil$ and $H = L - 1$. Then $L \leq 99/(\rho\eta)$, and (II.289.1) implies $p > 2H$. Reduction modulo p is injective on the integer interval $[-H, H]$. Consequently the box $[-H, H]^{32}$ meets the universal relation line exactly in $c\lambda$, $-H \leq c \leq H$.

Call an arrangement exceptional if an unintended vector in that box vanishes on its feature values. A nonzero polynomial of total degree at most two in 32 variables has at most $2p^{31}$ zeros. Pairing u with $-u$, which has the identical zero set, bounds the number by

$$D \leq ((2H + 1)^{32} - (2H + 1))p^{31} \leq 2(2L)^{32}p^{31}. \quad (\text{II.289.5})$$

The polynomial zero estimate follows by induction: a polynomial whose degree in the last variable is t has a leading coefficient of degree at most $D_0 - t$; at most $(D_0 - t)p^{r-2}$ base choices kill that coefficient, and every other base choice permits at most t roots. This gives the standard $D_0 p^{r-1}$ bound.

Every nonexceptional arrangement has distinct physical vertices. A coincidence of two labels would yield a feature relation with just one coefficient 1 and one coefficient -1 , which is not a multiple of λ . This observation includes the arrangements with $h = 0$.

II.289.2 Positive Fourier coefficients and repeated vertices

Define the normalized Fejér kernel

$$K_L(z) = \left| \frac{1}{L} \sum_{n=0}^{L-1} z^n \right|^2 = \sum_{|j| < L} c_j z^j, \quad c_j = \frac{L - |j|}{L^2}. \quad (\text{II.289.6})$$

On the unit circle $0 \leq K_L \leq 1$, the coefficients are nonnegative, $c_j = c_{-j}$, and $\sum_j c_j = 1$. Put

$$b_K = L^{-32}, \quad A_K = \sum_{|j| < L} c_j^{32} = b_K R_L, \quad R_L = 1 + 2 \sum_{j=1}^{L-1} (1 - j/L)^{32}.$$

Convexity of $(1 - x)^{32}$ and the trapezoidal bound for its integral on $[0, 1]$ give

$$R_L \geq \frac{2L}{33} \geq \frac{4}{\rho\eta}, \quad A_K \geq \frac{2}{33} L^{-31}. \quad (\text{II.289.7})$$

Choose independent uniform $t, u \in \mathbb{F}_p$. Conditional on these seeds, retain each physical point $v \in A$ independently with probability $K_L(e_p(t\phi(v) + uF(v)))$. For a respected arrangement,

expanding the product over its 32 labels and averaging over t, u retains all intended terms $c\lambda$. Their total contribution is A_K , since all other surviving coefficients are nonnegative.

This is also a lower bound for the *actual* survival probability when a physical point has several labels: the labelled product raises a weight in $[0, 1]$ to a higher power, whereas actual selection uses that weight once. Repeated vertices must be handled in this direction; independent selection of labels would describe a different random subset.

For a nonexceptional, unrespected arrangement the physical vertices are distinct, and every surviving feature relation is $c\lambda$. Its nonzero ϕ -defect forces $c = 0$, because $|c| \leq H < p$. Its survival probability is therefore exactly b_K . Each exceptional bad arrangement is bounded simply by one.

II.289.3 One random outcome has both required properties

Let X, Y be the numbers of respected and unrespected arrangements in the selected subset, and put $T_0 = \beta^{15}p^{32}$. Then

$$\mathbb{E}X \geq \rho A_K T_0, \quad \mathbb{E}Y \leq b_K T_0 + D.$$

The threshold (II.289.1), together with $L \leq 99/(\rho\eta)$, gives

$$D \leq 2(2L)^{32}p^{31} \leq L^{-32}\beta^{15}p^{32} = b_K T_0,$$

since $2 \cdot 99^2 = 18 \cdot 33^2$. Hence

$$\mathbb{E}(X - Y/\eta) \geq (\rho A_K - 2b_K/\eta)T_0 \geq \frac{\rho A_K}{2}T_0 > 0.$$

Choose an outcome with score at least its expectation. Then $Y \leq \eta X$ and

$$X \geq \frac{\rho A_K}{2}T_0 \geq \frac{\rho L^{-31}}{33}T_0 \geq \frac{\rho^{32}\eta^{31}}{3^{31}33^{32}}T_0.$$

These are (II.289.2)–(II.289.3), proving Theorem II.289.1. No independence between different arrangements is needed.

II.290 Improved quantitative propagation and the printed exponent

Source 57, delivered Section 4.

II.290.1 Keeping the exact initial density exponent

Proposition II.290.1 (Prepared frequency graph). *Let $0 < \alpha < 1$, put $a = \alpha/2$, and let $|f| \leq 1$ with $Q_4(f) > \alpha$. There is a set $A_2 \subseteq \mathbb{F}_p^2$ of density*

$$\beta \geq a^b, \quad b = 14\,409\,429, \tag{II.290.1}$$

and a separately order-eight Freiman map $\phi : A_2 \rightarrow \mathbb{F}_p$ such that $|\widehat{\Delta_h \Delta_k f}(\phi(h, k))| \geq a$ on A_2 .

Proof. For each (h, k) , write $M(h, k)$ for the largest Fourier modulus of $\Delta_h \Delta_k f$. By (II.288.1)–(II.288.2),

$$\alpha < Q_4(f) \leq \mathbb{E}_{h,k} M(h, k)^2,$$

because the sum of the squared Fourier moduli of each second derivative is at most one. Let $A_0 = \{(h, k) : M(h, k) \geq a\}$, of density d_0 . Then $\alpha < d_0 + (1 - d_0)a^2$. Since $\alpha = 2a$ and $0 < a \leq 1/2$, this implies $d_0 \geq a$. Choose a maximizing frequency at each pair in A_0 .

For each fixed h , let u_h be its fibre density. Apply Input theorem II.288.3 to $g = \Delta_h f$ and that fibre; empty fibres contribute zero and need no invocation. Summing the retained sizes and using convexity gives a set of density

$$\beta_1 \geq 2^{-1882} a^{9312} \mathbb{E}_h u_h^{1165} \geq 2^{-1882} a^{10477} \geq a^{12359}.$$

The last inequality uses $a \leq 1/2$. Repeat the same argument in the other coordinate. Restriction preserves the first Freiman property, and

$$\begin{aligned} \beta &\geq 2^{-1882} a^{9312} \beta_1^{1165} \\ &\geq a^{1882+9312+12359 \cdot 1165} = a^{14\,409\,429}. \end{aligned}$$

The Fourier lower bound survives both restrictions. □

The familiar rounded bound $\beta \geq a^{2^{24}}$ follows immediately. Retaining b instead of rounding it is useful: it lowers the final power-of-two density exponent by one more factor of two.

Proposition II.290.2 (Improved purification density). *Under the hypotheses of Proposition II.290.1, suppose $p \geq a^{-2^{37}}$. There is $A_3 \subseteq A_2$ with*

$$|A_3|/p^2 \geq \delta := a^d, \quad d = 2\,996\,201\,429, \quad (\text{II.290.2})$$

such that ϕ remains separately Freiman of order eight, every selected Fourier coefficient has modulus at least a , and

$$R(A_3, \phi) \geq (1 - 2^{-44})T(A_3).$$

Proof. Put $\beta_0 = a^b$, $\eta = 2^{-44}$ and $\rho_0 = \beta_0^{97} a^{336}$. For the actual density $\beta \geq \beta_0$, Input theorem II.288.4 yields

$$R(A_2, \phi) \geq \beta^{112} a^{336} p^{32} \geq \rho_0 \beta^{15} p^{32}.$$

Thus the lower proportion parameter is fixed independently of the actual density.

The right side of (II.289.1) is largest when $\beta = \beta_0$. Since $18 \cdot 33^2 < 2^{17}$, a sufficient upper bound for it is

$$\begin{aligned} 2\beta_0^{-15} \left(\frac{18 \cdot 33^2}{(\rho_0 \eta)^2} \right)^{32} &\leq a^{-(6223b+24865)} \\ &= a^{-89\,669\,901\,532} \leq a^{-2^{37}}. \end{aligned} \quad (\text{II.290.3})$$

Here $2^j \leq a^{-j}$ for every nonnegative integer j . The field-size and characteristic requirements of Theorem II.289.1 are both included in this bound.

Apply that theorem. With $C_0 = 2^{-1364}/(3^{31} 33^{32}) \geq 2^{-1618}$, its mass bound gives

$$\begin{aligned} R(A_3, \phi) &\geq C_0 \beta_0^{3104} \beta^{15} a^{10752} p^{32} \\ &\geq a^{3119b+10752+1618} p^{32} = a^{44\,943\,021\,421} p^{32}. \end{aligned}$$

Writing $\beta' = |A_3|/p^2$, Lemma II.288.2 implies $(\beta')^{15} \geq a^{44\,943\,021\,421}$. Therefore

$$\beta' \geq a^{44\,943\,021\,421/15} \geq a^d, \quad d = \left\lceil \frac{44\,943\,021\,421}{15} \right\rceil.$$

The ratio statement is (II.289.3), and the other properties are hereditary under restriction. $\square$

II.290.2 An explicit all-parameter length bound

Theorem II.290.3 (Quantitative extraction from the improved preparation). *For every $0 < \alpha < 1$, all sufficiently large primes p , and every $f : \mathbb{F}_p \rightarrow \mathbb{C}$ with $|f| \leq 1$ and $Q_4(f) > \alpha$, there are proper, equal-length progressions P, Q with the same nonzero step, a set $B \subseteq P \times Q$, and a formula (II.288.3), such that, writing $m = |P| = |Q|$,*

$$m \geq p^{2^{-(2/\alpha)^{2^{53}}}}, \quad (\text{II.290.4})$$

$$|B| \geq (\alpha/2)^{2^{41}} m^2, \quad (\text{II.290.5})$$

$$|\widehat{\Delta_h \Delta_k f}(\mu(h, k))| \geq \alpha/2 \quad ((h, k) \in B). \quad (\text{II.290.6})$$

More precisely, (II.290.4)–(II.290.5) may be replaced by $m \geq p^{\gamma(a^d)}$ and $|B| \geq 2^{-137} a^{704d} m^2$, with $a = \alpha/2$ and d as in (II.290.2).

Proof. Use Proposition II.290.2 and apply Input theorem II.288.5 at the lower density $\delta = a^d$. Its hypotheses use the actual density only through the lower bound δ , so the threshold is uniform in the function and all selected sets. The exact bounds follow directly from (II.288.4).

For the simpler powers of two, Lemma II.288.6 gives $\gamma(a^d) \geq 2^{-15Q(a^d)}$. Since $a \leq 1/2$,

$$\begin{aligned} 15Q(a^d) &= 15 \cdot 2^{2^{20}} a^{-d^{2^{21}}} \\ &\leq a^{-(d^{2^{21}} + 2^{20} + 4)} \leq a^{-2^{53}}, \end{aligned} \quad (\text{II.290.7})$$

where the penultimate exponent is $6\,283\,489\,820\,278\,788 < 2^{53}$. Also

$$2^{-137} a^{704d} \geq a^{137+704d} = a^{2\,109\,325\,806\,153} \geq a^{2^{41}}. \quad (\text{II.290.8})$$

Agreement with ϕ on B transfers the Fourier lower bound. $\square$

[write] **7 October 2026; sources 54–59.** *Checked at the write* (exact integers). $12359 \cdot 1165 + 1882 + 9312 = 14,409,429 = b$; $6223b + 24865 = 89,669,901,532 \leq 2^{37}$; $3119b + 10752 + 1618 = 44,943,021,421$ and $d = \lceil 44,943,021,421/15 \rceil = 2,996,201,429$; $d \cdot 2^{21} + 2^{20} + 4 = 6,283,489,820,278,788 < 2^{53}$; $137 + 704d = 2,109,325,806,153 < 2^{41}$. The comparison $(2/\alpha)^{2^{53}} \leq \alpha^{-2^{70}}$ is $(2^{17} - 1) \log(1/\alpha) \geq \log 2$, which holds exactly for $\alpha \leq e^{-\log 2/(2^{17}-1)} = 0.99999471$, so the endpoint range $1 - \alpha \leq 2^{-10}$ covers the rest with room. *Against the catalogue.* `theorem_13_12` (Sections12_13.lean, line 718) asks, for $0 < \alpha \leq 1$ with $\neg \text{UniformOfDegree } f \ \alpha \ 3$ (that is, $\|f\|_{U^4}^{16} > \alpha$), for proper `ModAPs` of common nonzero step and common length at least $N^{(1/2)^{(1/\alpha)^{2^{70}}}}$, a set B of at least $(\alpha/2)^{2^{76}}$ times the square, a map `BilinearOn` the whole square, and $\alpha N/2 \leq |\text{secondDifferenceFourier}|$ on B (unnormalized, with backward differences, whose moduli agree, as this source says); $\alpha = 1$ is vacuous. Theorem F.62.1 gives all of it with the stronger density $(\alpha/2)^{2^{41}}$. *The formal proof* (since `d48e8af00`; `theorem_13_12_holds`, `Proofs13EndpointPrintedTheorem.lean`, line 73) follows this route with the repository's own constants: a uniform Fejér purification with kernel parameter $\lceil 2^{34}/(\rho\eta) \rceil$ and retained mass $2^{-1150} \rho^{32} \eta^{31} \beta^{15} N^{32}$ (`section13_fejer_purification_uniform`,

`Proofs13FejerUniformPurification.lean`, line 10; here the mass constant is $1/(3^{31}33^{32})$, density $2^{-77}\rho^3\eta^3\beta$ after selection (`section13_fejer_density_uniform`, `Proofs13FejerPurifiedDensity.lean`, line 45), and with the same $a = \alpha/2$, $\delta = a^{14,409,429}$, $\rho = \delta^{97}a^{336}$, $\eta = 2^{-44}$ the graph density $a^{4,207,554,485}$ (`section13_fejer_frequency_graph`, `Proofs13FejerFrequencyGraph.lean`, line 42; $3(97 \cdot 14,409,429 + 336) + 14,409,429 + 209 = 4,207,554,485$, against d here). Hence the formal exponents are $4,207,554,485 \cdot 2^{21} + 2^{20} + 4 = 8,823,881,304,375,300 < 2^{53}$ for the length, as here, and $137 + 704 \cdot 4,207,554,485 = 2,962,118,357,577$, between 2^{41} and 2^{42} , for the density (`theorem_13_12_with_fejer_explicit_bound`, `Proofs13FejerExplicitExponent.lean`, line 55): the formal density exponent is 2^{42} , and this source's 2^{41} is not formalized. The printed length is reached for $\alpha \leq 1023/1024$ by Bernoulli's inequality (`theorem_13_12_away_from_one`, `Proofs13FejerPrintedRange.lean`, line 39) and for $\alpha > 1023/1024$ by this source's endpoint, $P = Q$ the whole group (`theorem_13_12_near_one`, `Proofs13EndpointPrintedTheorem.lean`, line 12). The repository has also proved the length exponent 2^{61} in place of 2^{70} , with density $(\alpha/2)^{2^{42}}$, for all $0 < \alpha \leq 1$ by the larger overlap $\alpha \leq 332/333$ (`theorem_13_12_with_improved_printed_bounds`, `Proofs13ImprovedPrintedBounds.lean`, line 11); that strengthening is the repository's, not this source's. `FORMALIZATION_STATUS.txt` credits source 57's route.

Remark II.290.4 (The size threshold). One may take

$$N_0(\alpha) = \max \left\{ 2, \left\lceil (\alpha/2)^{-2^{37}} \right\rceil, N_{\text{sq}}((\alpha/2)^d) \right\}.$$

This is a precise uniform threshold defined using the inherited extraction threshold. The power $(\alpha/2)^{-2^{37}}$ is an explicit sufficient threshold for *purification*; it is not asserted to bound the whole extraction. No closed elementary expression for N_{sq} is derived in this article.

[write] **7 October 2026; sources 54–59.** *The threshold made explicit* (checked at the write; this answers Research question V.216.8). Input theorem II.288.5 is the repository's `section13_complete_square_extraction_uniform_density` (`Proofs13UniformSquareExtraction.lean`, line 15). Source 53's Theorem F.58.1 proves its conclusion for every prime $p \geq N_0(\delta) = \lceil (8/c(\delta))^{1/\gamma(\delta)} \rceil$, $c = \zeta_{13}/2$, with side at least $p^{3\gamma(a)}$ and mass $2^{-137}a^{704}$ at the actual density $a \geq \delta$; since γ increases with the density (a^{1856} increases and $Q(a) = 2^{2^{20}}a^{-2^{21}}$ decreases), these are at least $p^{\gamma(\delta)}$ and $2^{-137}\delta^{704}$. So one may take $N_{\text{sq}}(\delta) = \lceil (8/c(\delta))^{1/\gamma(\delta)} \rceil$ and

$$N_0(\alpha) = \max \left\{ 2, \left\lceil (\alpha/2)^{-2^{37}} \right\rceil, \left\lceil (8/c(a^d))^{1/\gamma(a^d)} \right\rceil \right\}, \quad a = \alpha/2,$$

a closed threshold for Theorem F.62.1. Independently, the repository has made the threshold of the same interface explicit, `section13GeometricThreshold` (`Proofs13ExplicitGeometricThreshold.lean`, line 12) for `section13_complete_square_extraction_explicit` (line 24; 93d436414, 7 October 2026, after this source's pin), another formula. Neither was available to this source. Source 53 was written without knowledge of this source, and conversely.

II.290.3 Why an endpoint argument is necessary

The length exponent in (II.290.4) does not by itself dominate the printed exponent for every α approaching one. The factor $2/\alpha$ matters inside these powers. We retain that distinction rather than absorbing it informally.

For $\alpha \leq 1 - 2^{-10}$, comparison of the two exponents reduces to

$$(2^{17} - 1) \log(1/\alpha) \geq \log 2. \tag{II.290.9}$$

Indeed it is equivalent to $(2/\alpha)^{2^{53}} \leq \alpha^{-2^{70}}$. Now $\log(1/\alpha) \geq 1 - \alpha \geq 2^{-10}$, and $(2^{17} - 1)2^{-10} > 1 > \log 2$, so (II.290.9) holds.

For the remaining interval, an elementary rigidity argument gives a much stronger ambient square: Theorem II.247.5 proves that $Q_4(f) \geq 1 - \varepsilon$ implies

$$|B| \geq (1 - 332\varepsilon)p^2, \quad |\widehat{\Delta_h \Delta_k f}(chk)| \geq 3/4 - 4\varepsilon \quad ((h, k) \in B). \quad (\text{II.290.10})$$

With $\varepsilon = 1 - \alpha \leq 2^{-10}$, these bounds are at least $173p^2/256$ and $191/256$. Take $P = Q = \mathbb{F}_p$, both represented properly with step one. Their length p dominates the printed lower bound, the coefficient exceeds $\alpha/2$, and their good-pair density exceeds $a^{2^{41}}$.

Proof of Theorem F.62.1. Use Theorem II.290.3 and (II.290.9) when $0 < \alpha \leq 1 - 2^{-10}$. Use (II.290.10), proved in Theorem II.247.5, when $1 - 2^{-10} < \alpha < 1$. The threshold in Remark II.290.4 works for the combined statement. At $\alpha = 1$ the strict hypothesis $Q_4(f) > \alpha$ is impossible; the non-strict equality case is separately included in Theorem II.247.5 with $\varepsilon = 0$. $\square$

The new mechanism is the reduction of the density loss before the geometric extraction. A change from $a^{2^{66}}$ to a^d in that input is amplified by the factor $\delta^{-2^{21}}$ inside $\mathcal{Q}(\delta)$. The endpoint argument then removes the remaining parameter-range obstruction. Neither step changes an integer progression length into a real interval or assumes unrestricted versions of the raw catalogue extraction statements.

II.291 Source 57's chapter: bounded Schur deficits and the third energy level [write]

Added in the tenth write, 7 October 2026.

The next section is source 57's Section 6. *Logically it continues source 45's chapter* on the top of the torsion-free energy spectrum (Section II.183); it is printed here because no number printed earlier may move.

What is proved. For a positive real set C of n elements with Schur deficit $h(C) = s$ and $n \geq 5s - 1$, C is a positive dilate of $\{j + \lambda_j : 1 \leq j \leq n\}$ with (λ_j) nondecreasing, nonnegative and summing to s , so there are exactly $p(s)$ classes (Theorem II.292.3); for $s = 2$ the threshold $n = 6$ is optimal (Theorem II.292.4). For $m \geq 7$ the m -element torsion-free sets with $E(A) = M_m - 8(m - 3)$, $M_m = (2m^3 + m)/3$, are exactly the affine images of $J_m = \{0, \dots, m\} \setminus \{2\}$ and $K_m = \{0, 2, 3, \dots, m-1, m+1\}$; at $m = 6$ there is one further class, $\{0, 1, 2, 4, 5, 6\}$ (Theorem II.292.7). These answer source 45's first two research questions, its Questions 6.1–6.2 (Question V.166.1 and the question after it; item II.53); the energy recurrence and the first two levels are source 45's and Hegyvári's, credited. The results concern ordered real and torsion-free sets; the transfer to cyclic groups needs a Freiman isomorphism, as source 57 says.

Relation to the formal statements. No declaration concerns energy levels or Schur deficits; nothing in this chapter is formalized.

II.292 Bounded Schur defects and integer partitions

Source 57, delivered Section 6.

The ordered endpoint of Gowers's Section 7 has a discrete structure that is invisible in a bound with an unspecified constant: an arithmetic progression maximizes the additive energy of an m -element torsion-free set, and the next attainable energies have exact linear deficits. Source 45 of the

pinned repository [234] determined the first three values and classified the first two equality cases. It explicitly left open the classification at the third value and a general question about positive sets with a bounded deficit in their Schur-triple count. We answer both questions here.

The two conclusions have different scopes. Theorem II.292.3 classifies every fixed Schur deficit by integer partitions, beyond an explicit linear threshold in the deficit. Theorem II.292.7 classifies every third-level energy extremizer, with the optimal threshold of seven points. These statements concern real or torsion-free sets. They require an explicit rectification hypothesis before being applied to subsets of a finite cyclic group.

The maximum Schur count and its arithmetic-progression equality case are classical; see Huczynska [231, Theorem 9]. Her paper studies attainable counts in intervals, with its near-extremal structural section focused on the lower endpoint. The energy recurrence and congruence used below occur in Hegyvári [189, Section 2.1]. We reproduce short proofs of the required earlier ingredients to make the new classification self-contained. The additions relative to the pinned repository are the bounded-deficit partition theorem, the complete Schur-defect-two classification, and the third-energy equality theorem; no worldwide publication-priority assertion is intended.

II.292.1 Two consecutive perfect reflection rows

Let $C = \{c_1 < \cdots < c_n\} \subset (0, \infty)$, and define

$$S(C) = \#\{(x, y, z) \in C^3 : x + y = z\}, \quad h(C) = \binom{n}{2} - S(C). \quad (\text{II.292.1})$$

The summands are ordered and may coincide. Write r_j for the number of ordered representations $x + y = c_j$ with $x, y \in C$, and put

$$e_j = j - 1 - r_j. \quad (\text{II.292.2})$$

Both summands precede c_j , and each first summand determines the second. Consequently $e_j \geq 0$ and $h(C) = \sum_j e_j$. We call row j *perfect* when $e_j = 0$. Equivalently, its preceding prefix is invariant under the reflection $x \mapsto c_j - x$.

Lemma II.292.1 (Two perfect rows). *If rows j and $j + 1$ are perfect, then $c_i = id$ for $1 \leq i \leq j + 1$, where $d = c_1 > 0$.*

Proof. The case $j = 1$ is immediate. For $j \geq 2$, put

$$P = \{c_1, \dots, c_{j-1}\}, \quad Q = P \cup \{c_j\}, \quad t = c_{j+1} - c_j > 0.$$

Perfectness gives $c_j - P = P$ and $c_{j+1} - Q = Q$, whence

$$P + t = c_{j+1} - (c_j - P) \subseteq Q.$$

The left side has $j - 1$ points and omits $\min Q = c_1$, so it equals $Q \setminus \{c_1\}$. Comparing the increasing lists gives $c_{i+1} = c_i + t$ for $1 \leq i < j$. The first reflection also gives $c_j = c_1 + c_{j-1}$. Substitution yields $c_1 = t$, and therefore $c_i = it$ for $i \leq j + 1$. $\square$

In particular, after the first imperfect row there cannot be two consecutive perfect rows. This simple observation controls the entire finite pattern, rather than only the last representation count.

II.292.2 All fixed Schur deficits are eventually partitions

Theorem II.292.2 (A bounded terminal perturbation). *Let $s \geq 1$ be an integer. If $|C| = n \geq 3s + 1$ and $h(C) \leq s$, then for some $d > 0$ the set $d^{-1}C$ consists of positive integers and*

$$\{1, \dots, n - 2s\} \subseteq d^{-1}C \subseteq \{1, \dots, n + s\}. \quad (\text{II.292.3})$$

Proof. If every row is perfect, successive reflections give $C = d\{1, \dots, n\}$. Otherwise let k be the first imperfect row. Before it, $c_i = id$ with $d = c_1$. The suffix of rows $k, \dots, n$ starts with an imperfect row, has at most s imperfect rows, and has no two consecutive perfect rows by Lemma II.292.1. It therefore has at most $2s$ rows. Thus

$$L := k - 1 \geq n - 2s \geq s + 1.$$

For every $j \geq k$, $r_j = j - 1 - e_j \geq L - s \geq 1$. Hence c_j is a sum of two earlier points, and induction makes $a_j = c_j/d$ an integer.

For $j \geq k$, put $g_j = a_j - a_{j-1}$. Every $x \in \{1, \dots, \min(g_j - 1, L)\}$ is an earlier point, but $a_j - x > a_{j-1}$, so its partner is absent. Therefore

$$e_j \geq \min(g_j - 1, L).$$

Since $e_j \leq s < L$, we have $g_j - 1 < L$ and $e_j \geq g_j - 1$. Summing over $j \geq k$ gives

$$a_n - n = \sum_{j=k}^n (g_j - 1) \leq \sum_{j=k}^n e_j \leq s.$$

This proves both inclusions in (II.292.3). □

Theorem II.292.3 (Partition classification). *Let $s \geq 1$ and $n \geq 5s - 1$ be integers. A positive real n -set C satisfies $h(C) = s$ if and only if*

$$C = d\{j + \lambda_j : 1 \leq j \leq n\} \quad (\text{II.292.4})$$

for some $d > 0$ and integers

$$0 \leq \lambda_1 \leq \dots \leq \lambda_n, \quad \sum_{j=1}^n \lambda_j = s. \quad (\text{II.292.5})$$

There are exactly $p(s)$ such sets up to positive dilation, where $p(s)$ is the ordinary partition function. In every model, $\lambda_1 = \dots = \lambda_{n-s} = 0$.

Proof. Since $5s - 1 \geq 3s + 1$, Theorem II.292.2 applies. Normalize $d = 1$ and write $C = \{a_1 < \dots < a_n\}$ with integer entries. Let $H = \{1, \dots, a_n\} \setminus C$ be its holes. Every hole is at least $n - 2s + 1$, whereas $a_n \leq n + s$. The threshold gives

$$2(n - 2s + 1) > n + s \geq a_n. \quad (\text{II.292.6})$$

Thus no sum of two holes belongs to C .

For a point a_j , exactly $a_j - j$ integers below it are holes. For each such hole z , the positive integer $a_j - z$ cannot be another hole, by (II.292.6); hence it is an earlier point of C . It supplies one failed reflection entry in row j . Conversely, every failed entry arises in this way. Therefore

$$e_j = a_j - j. \quad (\text{II.292.7})$$

Set $\lambda_j = a_j - j$. Strict increase of the integer a_j makes these numbers nonnegative and nondecreasing; their sum is $h(C) = s$. This proves necessity.

Conversely, suppose (II.292.5) holds. There are at most s positive entries, so $a_j = j$ for $j \leq n - s$ and $a_n \leq n + s$. Every hole is at least $n - s + 1$, and $2(n - s + 1) > n + s$. Again no sum of two holes belongs to C , so the same argument proves (II.292.7) and $h(C) = s$.

Finally $n > s$, so $\lambda_1 = 0$ and the dilation is fixed by $d = c_1$. The positive tail of (λ_j) is exactly an integer partition of s , written in increasing order and padded by zeros. This gives the count $p(s)$. $\square$

[write] **7 October 2026; sources 60–69.** *Checked at the eleventh write* (Section II.296): source 66 proves the same normal form for $n \geq 3r + 1$ (its Theorem II.302.1), better than $5s - 1$ for $s \geq 3$; with Theorem II.292.4 and source 66's obstruction at $n = 2s + 1$, the optimal threshold satisfies $2s + 2 \leq N(s) \leq 3s + 1$ for $s \geq 3$, and source 66 conjectures $2s + 2$. The two sources were written the same day; neither cites the other.

For example, at deficit three the only eventual models are

$$\begin{aligned} &\{1, \dots, n - 1, n + 3\}, \\ &\{1, \dots, n - 2, n, n + 2\}, \\ &\{1, \dots, n - 3, n - 1, n, n + 1\}, \end{aligned}$$

up to positive dilation. The theorem supplies the sufficient threshold $n \geq 14$ in this example; it does not assert that this threshold is optimal. For deficit two we next obtain the optimal threshold directly.

II.292.3 The complete deficit-two classification

Theorem II.292.4 (Two missing Schur representations). *For $n \geq 6$, the positive real n -sets with $h(C) = 2$ are precisely*

$$d\{1, \dots, n - 1, n + 2\} \quad \text{and} \quad d\{1, \dots, n - 2, n, n + 1\}, \quad d > 0. \quad (\text{II.292.8})$$

At $n = 5$ there is exactly one additional dilation class, $d\{1, 2, 4, 5, 6\}$. Thus six is the optimal eventual threshold.

Proof. Let k be the first imperfect row and put $d = c_1$. The preceding points are $c_i = id$. There is either one row of defect two or two rows of defect one.

Suppose first that $k \geq 4$ and $e_k = 2$. Since $r_k = k - 3 > 0$, $c_k = qd$ for an integer q . For $k \leq q \leq 2k - 2$, the number of ordered pairs from $\{1, \dots, k - 1\}$ with sum q is $2k - 1 - q$; above this range it is zero. The condition $r_k = k - 3$ therefore gives $q = k + 2$. A following row would have to be perfect. Reflection of $\{1, \dots, k - 1, k + 2\}$ would have sum $k + 3$, sending the present point 2 to the absent point $k + 1$. Hence there is no following row, and $k = n$, giving the first model.

If $k \geq 4$ and $e_k = 1$, the same calculation gives $c_k = (k + 1)d$. The next row cannot be perfect: its reflection would send 2 to the missing k . Thus $e_{k+1} = 1$, and all later rows must be perfect. Normalize $d = 1$. The next point $q = c_{k+1}$ is an integer, because $r_{k+1} = k - 1 > 0$, and $q \geq k + 2$. Its representation count in $\{1, \dots, k - 1, k + 1\}$ is

$$\max(2k - 1 - q, 0) + 2\mathbf{1}_{\{1 \leq q - k - 1 \leq k - 1\}} + \mathbf{1}_{\{q = 2k + 2\}}.$$

At $q = k + 2$ this equals $k - 1$. At larger q it is at most $k - 2$: the interval contribution is at most $\max(k - 4, 0)$, the mixed contribution is at most two, and the isolated double-summand case has value one. Hence $q = k + 2$. A further perfect row would reflect $\{1, \dots, k - 1, k + 1, k + 2\}$ with sum $k + 3$, sending the present point 3 to the missing k . Thus $n = k + 1$, giving the second model.

For $k = 3$, the first two points are $d, 2d$. If $e_3 = 2$, the third row has no representations. A following perfect row would make the three-point prefix symmetric and force $c_3 = 3d$, a contradiction. If $e_3 = 1$, then $c_3 = 4d$. The next row must have defect one, and its two representations force $c_4 = 5d$ or $6d$. The prefix $\{1, 2, 4, 6\}$ is not symmetric, so the latter choice cannot continue. The prefix $\{1, 2, 4, 5\}$ is symmetric with sum six, allowing $c_5 = 6d$; then $\{1, 2, 4, 5, 6\}$ is not symmetric. This yields exactly the extra five-point model and no model of size six or more.

It remains to exclude $k = 2$ when $n \geq 5$. Write $c_1 = a, c_2 = b$, where $0 < a < b$ and $b \neq 2a$. Here $e_2 = 1$. If $e_3 = 0$, then $c_3 = a + b$. The fourth row cannot be perfect, because symmetry of $\{a, b, a + b\}$ would force $b = 2a$. Its defect must therefore be one, and its two representations give $c_4 = 2a + b$ or $a + 2b$. Neither four-point prefix is symmetric: comparison of its inner and outer pair sums would force respectively $b = 2a$ or $a = 0$. If instead $e_3 = 1$, its unique representation gives $c_3 = 2a$ or $2b$. A subsequent perfect row requires $a + c_3 = 2b$. The second option is impossible, while the first gives $b = 3a/2$ and a four-point prefix proportional to $\{2, 3, 4, 6\}$, which again is not symmetric. Thus no case with $k = 2$ reaches five points.

The two models in (II.292.8) have final row defects (2) and (1, 1) respectively. The exceptional five-point model has row defects (0, 0, 1, 1, 0). This also verifies sufficiency. $\square$

II.292.4 Ordered energy and its first gaps

For a nonempty real set A of size m , define

$$E(A) = \#\{(a, b, c, d) \in A^4 : a + b = c + d\}, \quad M_m = \frac{2m^3 + m}{3}, \quad D(A) = \frac{M_m - E(A)}{4}. \quad (\text{II.292.9})$$

For B of size n and $a > \max B$, set

$$T(B, a) = \#\{(x, y, z) \in B^3 : x + y - z = a\}, \quad e(B, a) = \binom{n}{2} - T(B, a).$$

Each pair $z < x$ in B determines the only possible completion $y = a - x + z$. Thus $e(B, a)$ is a nonnegative integer. Counting the occurrences of the largest point in an additive quadruple gives

$$E(B \cup \{a\}) = E(B) + 4n + 1 + 4T(B, a),$$

and therefore

$$D(B \cup \{a\}) = D(B) + e(B, a). \quad (\text{II.292.10})$$

Iterating from a singleton also gives the useful exact identity

$$D(A) = \#\{x < y < z \text{ in } A : x + z - y \notin A\}. \quad (\text{II.292.11})$$

These identities prove that $D(A)$ is a nonnegative integer.

The transformation $C = a - B$ identifies $e(B, a) = h(C)$. Consequently, Theorem II.292.3 already has a direct endpoint interpretation: if $e(B, a) = s$ and $|B| \geq 5s - 1$, then $B \cup \{a\}$ is an affine image of

$$\{0\} \cup \{j + \lambda_j : 1 \leq j \leq |B|\},$$

with (λ_j) a partition of s as above. It is contained in an arithmetic progression of length at most $|B| + 1 + s$ and omits at most s points from it.

We recall the short endpoint facts needed for energy equality.

Lemma II.292.5 (Endpoint costs zero and one). *For $a > \max B$, one has $e(B, a) = 0$ if and only if $B \cup \{a\}$ is an arithmetic progression. If $|B| = n \geq 4$ and $e(B, a) = 1$, then, after translating $\min B$ to zero, for some $d > 0$,*

$$B = d\{0, 2, 3, \dots, n\}, \quad a = (n + 1)d. \quad (\text{II.292.12})$$

Proof. The zero assertion follows from the zero-defect Schur classification: $h(C) = 0$ forces $C = d\{1, \dots, n\}$ by its perfect rows. For defect one there is a unique imperfect row k . Before it $c_i = id$. If $k \geq 3$, its $k - 2 > 0$ representations force $c_k = (k + 1)d$. A following perfect row would reflect $2d$ to the missing kd , so $k = n$. If $k = 2$, perfect rows three and four would force first $c_3 = d + c_2$ and then $2c_2 = d + c_3$, contradicting $c_2 \neq 2d$. Thus $C = d\{1, \dots, n - 1, n + 1\}$. Since $a = \max C$ after $\min B = 0$, this is exactly (II.292.12). $\square$

For an AP $B = d\{0, \dots, n - 1\}$, a nonintegral value of a/d has $T(B, a) = 0$. For an integer $a/d = q \geq n$,

$$T(B, a) = \sum_{h=\max(1, q-n+1)}^{n-1} (n - h). \quad (\text{II.292.13})$$

Indeed a positive difference hd has $n - h$ representations, and its completion belongs to B precisely when $q - h \leq n - 1$. The successive exceptional endpoints $q = n, n + 1$ have costs $0, n - 1$; for $n \geq 4$, every other endpoint has cost at least $2n - 3$.

Define

$$H_m = \{0, 1, \dots, m - 2, m\}, \quad L_m = \{0, 2, 3, \dots, m\} = m - H_m.$$

An affine model permits both positive and negative nonzero slopes.

Proposition II.292.6 (The preceding energy gaps). *The following assertions hold for real sets.*

1. $D(A) = 0$ exactly for arithmetic progressions.
2. For $m \geq 3$, every non-AP m -set has $D(A) \geq m - 2$.
3. For $m \geq 5$, equality $D(A) = m - 2$ holds exactly for affine images of H_m .
4. For $m \geq 5$, outside the preceding two classes one has $D(A) \geq 2m - 6$.

Proof. The first statement follows from (II.292.10) and Lemma II.292.5. For the second, induct on m , with the three-point case given by (II.292.13). If removing the maximum leaves an AP, that same formula gives endpoint cost at least $m - 2$. Otherwise its defect is at least $m - 3$, and the endpoint cost is at least one. Equality is attained by H_m .

For the third assertion, an AP prefix forces the displaced endpoint in H_m . A non-AP prefix forces endpoint cost one, and Lemma II.292.5 gives L_m .

For the fourth assertion, the case $m = 5$ follows from integrality and the strict inequality $D(A) > 3$ outside the first two classes. Induct with $n = m - 1$ and $B = A \setminus \{a\}$, where $a = \max A$. An AP prefix has cost at least $2n - 3 > 2n - 4$ by (II.292.13). If B is second-level, normalize it as

H_n or L_n . Its defect is $n - 2$, and difference one has multiplicity $n - 2$. Noninteger endpoints have cost $\binom{n}{2}$; integer endpoints at least $n + 2$ fail on every difference-one pair. At $a = n + 1$, L_n extends to another excluded second-level set, whereas H_n fails on its $n - 2$ pairs of difference two. Hence all admissible endpoints cost at least $n - 2$. Finally, if B is in neither class, induction gives $D(B) \geq 2n - 6$ and endpoint costs zero and one are excluded by Lemma II.292.5. Its cost is at least two, giving $D(A) \geq 2n - 4 = 2m - 6$. $\square$

II.292.5 Every third-level extremizer

Theorem II.292.7 (Complete third-energy classification). *Let $m \geq 7$. For every m -element real set A ,*

$$E(A) = \frac{2m^3 + m}{3} - 8(m - 3) \quad (\text{II.292.14})$$

if and only if A is an affine image of one of

$$J_m = \{0, 1, \dots, m\} \setminus \{2\}, \quad K_m = \{0, 2, 3, \dots, m - 1, m + 1\}. \quad (\text{II.292.15})$$

At $m = 6$ the only additional affine class is $\{0, 1, 2, 4, 5, 6\}$. Thus seven is the smallest eventual threshold for the two-class statement. These classifications also hold in every torsion-free abelian group.

Proof. Write $n = m - 1$, $a = \max A$, and $B = A \setminus \{a\}$. The target defect is $D(A) = 2n - 4$. If B is an AP, the two exceptional endpoints put A in the first two energy classes, while every other endpoint costs at least $2n - 3$. Hence this case is impossible.

Suppose next that B is second-level. Since $n \geq 5$, positively normalize it as H_n or L_n . Its defect is $n - 2$, so its endpoint cost must be exactly $n - 2$. Noninteger endpoints cost $\binom{n}{2} > n - 2$. For H_n , the endpoint $n + 1$ fails exactly on difference two, whose multiplicity is $n - 2$. It produces a reflected J_m . The endpoint $n + 2$ fails on difference one with multiplicity $n - 2$ and also on difference three with positive multiplicity $n - 3$. Every endpoint at least $n + 3$ fails on both differences one and two, each of multiplicity $n - 2$. Thus no other endpoint qualifies.

For L_n , the endpoint $n + 1$ gives another second-level set and is excluded. The endpoint $n + 2$ fails exactly on difference one, of multiplicity $n - 2$, giving K_m . Every endpoint at least $n + 3$ fails on differences one and two, again each of multiplicity $n - 2$. This exhausts the second-level prefixes.

If B is in neither preceding class, Proposition II.292.6 gives $D(B) \geq 2n - 6$. Endpoint costs zero and one are excluded, so equality forces

$$D(B) = 2n - 6, \quad e(B, a) = 2.$$

Apply Theorem II.292.4 to $C = a - B$; here $n \geq 6$. After translating $\min B = 0$, its first model gives

$$B = d\{0, 3, 4, \dots, n + 1\}, \quad a = (n + 2)d.$$

Reflection turns this B into $d\{0, 1, \dots, n - 2, n + 1\}$. Formula (II.292.13) gives $D(B) = 2n - 5$, a contradiction. The second model gives

$$B = d(\{0, 1, \dots, n\} \setminus \{2\}), \quad a = (n + 1)d,$$

and hence $A = dJ_m$. This proves necessity.

For sufficiency, the ordered-completion identity gives, for $1 \leq j \leq m-1$,

$$D(\{0, 1, \dots, m\} \setminus \{j\}) = j(m-j) - \min(j, m-j). \quad (\text{II.292.16})$$

Indeed a failed completion must equal j . Its extremal entries satisfy $x < j < z$ and its middle entry is $y = x + z - j$. Of the $j(m-j)$ choices of (x, z) , precisely $\min(j, m-j)$ have $y = j$ and must be excluded. Taking $j = 2$ gives $D(J_m) = 2m - 6$. For K_m , its prefix L_{m-1} has defect $m - 3$, and its endpoint has exactly $m - 3$ failing pairs of difference one. Hence $D(K_m) = 2m - 6$.

For $m = 6$, the preceding prefix analysis is unchanged except that $|C| = 5$. Its extra Schur model gives $A = d\{0, 1, 2, 4, 5, 6\}$. Formula (II.292.16) gives its defect $6 = 2m - 6$, and Theorem II.292.4 supplies no further exceptions. Its normalized successive gaps are $(1, 1, 2, 1, 1)$, whereas those of J_6 are $(1, 2, 1, 1, 1)$ up to reversal and those of K_6 are $(2, 1, 1, 1, 2)$. It is therefore a distinct affine class, proving optimality of the threshold.

Finally the subgroup generated by a finite set in a torsion-free abelian group is isomorphic to $\mathbb{Z}^r$ and embeds additively into $\mathbb{R}$ by choosing r rationally independent real numbers. This preserves energy. Every classified model contains a pair separated by its unit step, so the real step is the image of an actual group difference. Injectivity pulls back every model relation, completing the torsion-free transfer. $\square$

[write] 7 October 2026; sources 54–59. *Source 45's first two research questions are answered* (Section II.183; item II.53): its Question V.166.1 by this theorem, with the optimal threshold $m \geq 7$ and the exception $\{0, 1, 2, 4, 5, 6\}$ at $m = 6$, and its question on bounded Schur defect by Theorems II.292.3 and II.292.4 (sufficient threshold $5s - 1$; optimal values 4, 6, 8 for $s = 1, 2, 3$ are not proved in general). The finite enumerations of this source were rerun on a copy at the write with identical output. Source 66, placed after this source (06b30ca52) and not yet written in, records the same partition normal form with threshold $3r + 1$, a third-level classification from $m \geq 8$ and the fourth level (from its placement record, not checked here).

The partition theorem also gives a finite route to further energy classifications: a bounded endpoint cost eventually permits only $p(s)$ positive dilation classes. It does not, by itself, prove that a bound $D(A) \leq K|A|$ always provides a bounded-cost endpoint of sufficient length. That additional localization is the next structural obstacle.

II.293 Source 67's chapter: affine rigidity in characteristic three [write]

Added in the eleventh write, 7 October 2026.

The next two sections are source 67's Sections 3–4. *Logically they continue the affine-rigidity chapters* of source 11 (Section II.22) and of sources 59 and 56 (Section II.248); they are printed here because no number printed earlier may move.

What is proved. For $f : G \rightarrow W$, G finite abelian of order N and W an elementary abelian 3-group, at error density t from an affine reference: $\varepsilon(f) \geq \Psi(t) = 4t - 10t^2 + 6t^3$ for every $t \in [0, 1]$, with a nonnegative remainder $(2 - 3t)(2a - c) + \frac{4}{3}b$ that measures several error colours and the target rank, and, for $0 < t < 2/3$, equality exactly for one nonzero error value on a subgroup coset; the endpoint $t = 2/3$ has its own affine equality family (Theorem II.294.1). Source 59 proved $\varepsilon \geq \Psi(t)$ on $t < 4/9$ for every target without involutions (Theorem II.250.4); elementary abelian 3-groups have no involutions but have 3-torsion, so source 67 extends that range to $t \leq 2/3$ for this target

class. With a Fourier argument: a unique nearest affine map, within $D(\varepsilon)$, whenever $\varepsilon < 1/4$ on arbitrary domains or $\varepsilon < 4/9$ on elementary abelian 3-domains, both thresholds sharp as universal strict thresholds (Theorem II.295.3); and the non-affine maximum $\mathcal{E}(f) \leq N^3 - 4N^2 + 10N - 6$ from $N \geq 13$, with 13 optimal (Corollary II.295.5). *Two different “4/9”*: source 59’s bounds the error density t , source 67’s the defect ε . The recovery thresholds compare with source 11’s $\varepsilon < 1/9$ and source 59’s $2/9$ (for all involution-free targets, through classical correction). The order-thirteen cutoff settles source 11’s cutoff question for elementary abelian 3-group targets; its sharpness example, the parity map $\mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/3\mathbb{Z}$ of energy $1296 > 1266$, is source 59’s (Section II.252), uncited at that point (note after the corollary). The thresholds concern the four-point graph-energy test only, as source 67 says. Sources 93 and 96, placed during this write, record the recovery threshold $1/4$ and the cutoff 13 for every abelian target (93) and the whole cubic range for every odd-order target (96); they are not checked here (note after Theorem II.294.1).

Relation to the formal statements (checked at HEAD fb9602e55). `proposition_6_1` and `lemma_6_2` (Sections06_07.lean, lines 37 and 47) have the exact companions `proposition_6_1_holds` and `lemma_6_2_holds` (Proofs05_10.lean, lines 505 and 589); the affine-rigidity theory has no formal counterpart, and nothing in this chapter is formalized.

II.294 Cubic rigidity throughout the nontrivial range in characteristic three

Source 67, delivered Section 3.

The cubic error profile in the preceding repository work is

$$\Psi(t) = 4t - 10t^2 + 6t^3 = 2t(1 - t)(2 - 3t).$$

Its nontrivial range is $0 < t < 2/3$. The result below establishes this entire range when the error values belong to an elementary abelian 3-group. The domain can be any finite abelian group. A second result gives exact universal thresholds for unique affine recovery: graph defect below $1/4$ on arbitrary finite abelian domains, and below $4/9$ on elementary abelian 3-domains. These are strict thresholds, and explicit examples show that neither endpoint can be included in the respective universal recovery assertion.

The target restriction is essential to the proof. Random linear projections to $\mathbb{F}_3$ retain enough of the signed graph energy to reduce the main inequality to a two-color autocorrelation estimate. This is more information than a pointwise maximum of a signed four-position kernel retains.

II.294.1 Definitions and the principal local theorem

Let G be a finite abelian group of cardinality N , let W be an elementary abelian 3-group, and let $f : G \rightarrow W$. The group W may be infinite: only the finite-dimensional vector space generated by the finitely many values used below is relevant. Define the graph energy and its normalized defect by

$$\begin{aligned} \mathcal{E}(f) &= \#\{(x_1, x_2, x_3, x_4) \in G^4 : x_1 + x_2 = x_3 + x_4, f(x_1) + f(x_2) = f(x_3) + f(x_4)\}, \\ \varepsilon(f) &= 1 - \frac{\mathcal{E}(f)}{N^3}. \end{aligned}$$

An affine map means a homomorphism plus a constant. Fix any affine reference ℓ , put $g = f - \ell$, and write

$$S = \{x : g(x) \neq 0\}, \quad s = |S|, \quad t = s/N, \quad c_v = |\{x : g(x) = v\}| \quad (v \neq 0).$$

The reference need not be a nearest affine map. Set

$$\begin{aligned} M &= \sum_{v \neq 0} c_v^2, & A &= s^2 - M, & C &= \sum_{v \neq 0} c_v c_{-v}, \\ Z &= \sum_{\substack{u, v \neq 0 \\ u+v \neq 0}} c_u c_v c_{u+v}, & T &= \sum_{\substack{u, v, w \neq 0 \\ \dim_{\mathbb{F}_3} \langle u, v, w \rangle \geq 2}} c_u c_v c_w. \end{aligned}$$

Thus A counts ordered unequal-color pairs, C counts ordered opposite-color pairs, Z counts weighted Schur triples, and T counts ordered triples not contained in one one-dimensional target subspace. Let $a = A/N^2$, $c = C/N^2$, and $b = T/N^3$. Always $0 \leq C \leq A$.

Theorem II.294.1 (Characteristic-three cubic rigidity). *With the preceding definitions, the following inequality holds for every $0 \leq t \leq 1$:*

$$\varepsilon(f) - \Psi(t) \geq (2 - 3t)(2a - c) + \frac{4}{3}b. \quad (\text{II.294.1})$$

Consequently $\varepsilon(f) \geq \Psi(t)$ for every $t \in [0, 1]$. For $0 < t < 2/3$, equality holds if and only if

$$f(x) = \ell(x) + v 1_{x_0+H}(x) \quad \text{for some } v \neq 0 \text{ and subgroup } H \leq G. \quad (\text{II.294.2})$$

At $t = 2/3$, equality holds if and only if $g = f - \ell$ is affine and has an image of cardinality three containing zero. Equivalently,

$$g(x) = v(\lambda(x) + j) \quad (v \neq 0, j \in \mathbb{F}_3),$$

where $\lambda : G \rightarrow \mathbb{F}_3$ is a surjective homomorphism.

[write] **7 October 2026; sources 60–69.** *Forward pointers* (from placement records, not checked at this write). Source 96, placed by 09f5b745a during this write, records $\varepsilon(f) \geq \Psi(t)$ on the whole range $t \leq 2/3$, with the coset equality classification, for *every* odd-order target, which would contain this theorem's inequality; source 93 records affine recovery below graph defect $1/4$ and the non-affine maximum from $N \geq 13$ for every abelian target (Theorem II.295.3 and Corollary II.295.5 here concern elementary abelian 3-group targets). Both are for the writes that add them.

The inequality for $2/3 < t < 1$ follows simply from $\Psi(t) < 0$ and $\varepsilon(f) \geq 0$. The mathematical improvement is the whole positive interval, the explicit color remainder in (II.294.1), and its rigidity statement. In particular, no failure of the cubic inequality beyond $2/3$ is being asserted.

II.294.2 The two-color autocorrelation estimate

For a finitely supported complex function h on an abelian group put

$$\mathcal{E}(h) = \sum_d \left| \sum_x h(x+d) \overline{h(x)} \right|^2.$$

For real h this equals the sum of $h(x_1)h(x_2)h(x_3)h(x_4)$ over additive quadruples.

Lemma II.294.2 (Two colors separated by an angle of $\pi/3$). *Let P, Q be disjoint finite sets in an abelian group, of cardinalities x, y , respectively. Put $s = x + y$ and $h = 1_P + e^{i\pi/3}1_Q$. Then*

$$\mathcal{E}(h) \leq s^3 - 2xy \max\{x, y\} \leq s^3 - sxy. \quad (\text{II.294.3})$$

If $x, y > 0$, equality in the second upper bound $\mathcal{E}(h) \leq s^3 - sxy$ holds precisely when $x = y$ and P, Q are distinct cosets of the same finite subgroup H , with their coset difference not of order two in the quotient group.

Proof. For a fixed direction d , write the autocorrelation as $A_d + e^{i\pi/3}B_d + e^{-i\pi/3}C_d$, where A_d counts the same-color pairs and B_d, C_d count the two orientations of cross-color pairs. Write $v = B_d + C_d$, $r = A_d + v$, and $m = \min\{x, y\}$. The elementary constraints are

$$A_d \geq 0, \quad 0 \leq B_d, C_d \leq m, \quad r \leq s.$$

Expansion of the squared modulus gives

$$\left| A_d + e^{i\pi/3}B_d + e^{-i\pi/3}C_d \right|^2 = r^2 - A_d v - 3B_d C_d.$$

We claim

$$r(s - r) + A_d v + 3B_d C_d \geq (s - m)v. \quad (\text{II.294.4})$$

After subtraction of the right side, the left side is

$$A_d(s - A_d - v) + v(m - v) + 3B_d C_d.$$

If $v \leq m$, every displayed term is nonnegative. If $v \geq m$, then $(m - B_d)(m - C_d) \geq 0$ gives $B_d C_d \geq m(v - m)$, and hence

$$v(m - v) + 3B_d C_d \geq (v - m)(3m - v) \geq 0,$$

since $v \leq 2m$. This proves (II.294.4). Now sum over d , using $\sum_d r = s^2$ and $\sum_d (B_d + C_d) = 2xy$. The first bound in (II.294.3) follows; the second uses $\max\{x, y\} \geq s/2$.

For equality in the second bound with $xy > 0$, first $x = y = m = s/2$. The proof of (II.294.4) shows that a direction with a cross-color pair has $v = m$, $B_d C_d = 0$, and $A_d \in \{0, m\}$. The possibility $A_d = m$ would give $r = s$, so translation by d would permute $P \cup Q$. A permutation of the union moves equally many points from P to Q and from Q to P , contradicting $B_d C_d = 0$ and $v > 0$. Therefore $A_d = 0$. For a direction with no cross-color pair, equality gives $A_d = 0$ or $A_d = s$. Consequently every positive autocorrelation of P has value $|P|$ and the same direction preserves Q . The usual equality case of $E(P) \leq |P|^3$ makes P a coset of a subgroup H ; the preceding preservation and $|Q| = |H|$ make Q another coset of H . The two oriented cross differences must be distinct, excluding order two for their difference. Conversely, these cosets have exactly the autocorrelation patterns just described, so they attain equality. $\square$

Lemma II.294.3 (Signed graph energy with two nonzero ternary values). *Suppose $g : G \rightarrow \mathbb{F}_3$ takes the values 1 and -1 on disjoint sets P, Q and is zero elsewhere. Define*

$$D(x, w) = 1_{\{g(x) \neq 0\}}(1_{\{g(x)\}}(w) - 1_{\{0\}}(w)).$$

If $|P| = x$, $|Q| = y$, and $s = x + y$, then

$$\mathcal{E}(D) \leq 6s^3 - 6sxy. \quad (\text{II.294.5})$$

When $xy > 0$, equality has exactly the coset description in Lemma II.294.2. More strongly, the term $6sxy$ can be replaced by $12xy \max\{x, y\}$.

Proof. Let $\omega = e^{2\pi i/3}$. Fourier expansion only in the target coordinate gives zero for its trivial character. For either other character, the domain function is $(\omega - 1)1_P + (\omega^2 - 1)1_Q$, or its conjugate. The two coefficients have modulus $\sqrt{3}$ and ratio $e^{i\pi/3}$ or its conjugate. Target orthogonality therefore gives

$$\mathcal{E}(D) = 6\mathcal{E}(1_P + e^{i\pi/3}1_Q).$$

Apply Lemma II.294.2. □

For a scalar target, the stronger version and the signed identity below already give the entirely explicit estimate

$$\varepsilon(f) - \Psi(t) \geq (2 - 6 \min\{|P|, |Q|\}/N) \frac{2|P||Q|}{N^2}. \quad (\text{II.294.6})$$

Thus an imbalance between the two nonzero error colors improves the color penalty as well.

II.294.3 Linear projection and the seven target configurations

The passage to vector-valued targets uses an exact projection identity, not a union bound over coordinates. Let V be the finite-dimensional subspace of W spanned by the error values, and let λ range uniformly over $V^* = \text{Hom}(V, \mathbb{F}_3)$, including zero. On $G \times V$, define

$$D = \sum_{x \in S} (\delta_{(x, g(x))} - \delta_{(x, 0)}).$$

Write D_λ for the pushforward of D under $(x, w) \mapsto (x, \lambda(w))$; the two entries cancel when $\lambda(g(x)) = 0$.

Lemma II.294.4 (Exact projected signed energy). *One has*

$$\mathcal{E}(D) = \frac{3}{2} \mathbb{E}_\lambda \mathcal{E}(D_\lambda). \quad (\text{II.294.7})$$

Proof. For every $v \in V$, $\Pr_\lambda[\lambda(v) = 0] = 1/3 + (2/3)1_{\{v=0\}}$. Apply this to the signed target sum of each horizontal additive quadruple in the expansion of $\mathcal{E}(D_\lambda)$. The constant $1/3$ part vanishes: for every x , $\sum_w D(x, w) = 0$, so the unrestricted sum over its four target coordinates factors into four zero sums. The remaining part is $(2/3)\mathcal{E}(D)$. □

For $a_1, a_2, a_3 \in \mathbb{F}_3$, define $\kappa(a_1, a_2, a_3)$ to be zero if any entry is zero, one if all three entries are equal and nonzero, and $2/3$ otherwise. If x_λ, y_λ are the two nonzero projected color sizes and $s_\lambda = x_\lambda + y_\lambda$, then

$$\sum_{u, v, w \neq 0} c_u c_v c_w \kappa(\lambda u, \lambda v, \lambda w) = s_\lambda^3 - s_\lambda x_\lambda y_\lambda.$$

Indeed, among three nonzero projected labels the number of ordered non-monochromatic triples is $3x_\lambda y_\lambda s_\lambda$. Lemmas II.294.3 and II.294.4 now imply

$$\mathcal{E}(D) \leq \sum_{u, v, w \neq 0} c_u c_v c_w K(u, v, w), \quad K(u, v, w) = 9\mathbb{E}_\lambda \kappa(\lambda u, \lambda v, \lambda w). \quad (\text{II.294.8})$$

For a triple of nonzero vectors, let U count its unequal pairs, P_- its opposite pairs, and R the equations among $u + v = w$, $u + w = v$, and $v + w = u$. At most one of these equations holds. The following table is exhaustive. In its rank-two rows, “independent” means that the displayed exceptional vector is independent of the repeated or opposite pair.

Configuration	U	P_-	R	K	$6 - (2U - P_- + 4R)/3 - K$
All three equal	0	0	0	6	0
One line, not all equal	2	2	1	4	0
Two equal, independent third	2	0	0	10/3	4/3
Opposite pair, independent third	3	1	0	8/3	5/3
Rank two, pairwise nonparallel, $R = 1$	3	0	1	4/3	4/3
Rank two, pairwise nonparallel, $u + v + w = 0$	3	0	0	2	2
Rank three	3	0	0	2	2

Here is a direct way to verify the numerical column without appealing to computation. In rank one, the nonzero functional values occur with probability $2/3$. In rank two, the values on a basis are the nine equally likely pairs in $\mathbb{F}_3^2$. The third row has two nonzero monochromatic outcomes and two nonzero non-monochromatic outcomes; the fourth has four of the latter. For $w = u + v$ there are two nonzero outcomes, both non-monochromatic. For $w = -u - v$ there are two nonzero outcomes, both monochromatic. In rank three, among the 27 outcomes there are two nonzero monochromatic and six nonzero non-monochromatic outcomes. Multiplying the averages by nine gives the table. The rank-two classification is exhaustive because any dependence of three pairwise nonparallel vectors over $\mathbb{F}_3$ has three nonzero coefficients, each equal to 1 or -1 .

Proposition II.294.5 (A signed energy bound retaining target rank). *With the notation of Theorem II.294.1,*

$$\mathcal{E}(D) \leq 6s^3 - 2sA + sC - 4Z - \frac{4}{3}T. \quad (\text{II.294.9})$$

Proof. The table gives

$$K(u, v, w) \leq 6 - \frac{2U - P_- + 4R}{3} - \frac{4}{3} 1_{\{\dim\langle u, v, w \rangle \geq 2\}}.$$

Sum with weight $c_u c_v c_w$. The sums of U , P_- , and R are $3sA$, $3sC$, and $3Z$, respectively. Equation (II.294.8) proves the assertion. $\square$

II.294.4 The exact graph expansion and equality

For completeness, the signed identity used to finish the argument is included here. Let $B = 1_{G \times \{0\}}$, so that $B + D$ is the graph indicator of g . Expansion in the number of D factors gives

$$\mathcal{E}(g) = N^3 - 4sN^2 + N(6s^2 + 4M + 2C) - 4s^3 - 8sM - 4sC + 4Z + \mathcal{E}(D). \quad (\text{II.294.10})$$

The one-factor contribution is $-4sN^2$. For two factors, four opposite-sign position pairs contribute $N(s^2 + M)$ each and two same-sign pairs contribute $N(s^2 + C)$ each. For three factors, put $m_0 = -s$, $m_v = c_v$ for $v \neq 0$; the target sum is

$$\sum_{u+v=w} m_u m_v m_w = Z - s^3 - 2sM - sC,$$

and there are four choices of the omitted position. The final term has four D factors. All these counts include repeated domain coordinates with their correct multiplicities.

Proof of Theorem II.294.1. Affine subtraction preserves graph energy. Subtracting the cubic profile from (II.294.10) gives

$$N^3(\varepsilon(f) - \Psi(t)) = 2(N - 2s)(2A - C) - 4Z + 6s^3 - \mathcal{E}(D).$$

Insert Proposition II.294.5 to obtain

$$N^3(\varepsilon(f) - \Psi(t)) \geq (2N - 3s)(2A - C) + \frac{4}{3}T.$$

This is (II.294.1). If $0 < t < 2/3$, its right side is nonnegative and its color term is at least $(2N - 3s)A$. Equality therefore forces $A = 0$, so all nonzero error values equal one v . For such an error, direct substitution or support-pattern counting gives

$$\varepsilon(v1_S) = \Psi(t) + \frac{6(s^3 - E(S))}{N^3}. \quad (\text{II.294.11})$$

The elementary inequality $E(S) \leq s^3$ is an equality precisely for subgroup cosets. Indeed, if $r_S(d) = |S \cap (S + d)|$, then $E(S) = \sum_d r_S(d)^2 \leq s \sum_d r_S(d) = s^3$; equality makes each positive $r_S(d)$ equal s , and the translation stabilizer of S then makes S a coset. This proves both directions of (II.294.2).

At $t = 2/3$, the profile is zero. Full graph energy is equivalent to affineness: the identity

$$\mathcal{E}(g) = \sum_h \sum_w |\{x : g(x+h) - g(x) = w\}|^2$$

shows that full energy is equivalent to every derivative being constant, and the derivative cocycle makes these constants a homomorphism. If an affine map g has a zero fiber of size $N/3$, its image is a three-element subgroup containing zero. This gives exactly the endpoint description in the theorem. $\square$

Proposition II.294.6 (Attainment of the color coefficient). *Let $H \leq G$, let P, Q be distinct H -cosets whose difference has order other than two in G/H , and let $g = v1_P - v1_Q$ for $v \neq 0$. Put $h = |H|$ and $t = 2h/N$. Then $A = C = 2h^2$, $T = 0$, and*

$$\varepsilon(g) - \Psi(t) = (2 - 3t) \frac{2h^2}{N^2}. \quad (\text{II.294.12})$$

Thus the color coefficient in (II.294.1) cannot be uniformly increased with the other terms fixed at these parameters.

Proof. Lemma II.294.3 has equality for these two cosets, and their two nonzero values lie on one target line. The projection table also has equality on that line. Consequently every inequality used in deriving (II.294.1) is an equality. $\square$

For example, take $G = \mathbb{Z}/q\mathbb{Z}$ and $H = \{0\}$ with $q \geq 3$. At $q = 3$ the two-color example is an affine map and gives the endpoint $t = 2/3$. At arbitrarily large q it gives genuine two-error examples and the exact color penalty at $t = 2/q$.

Corollary II.294.7 (Color stability). *If $0 < t < 2/3$ and $\varepsilon(f) \leq \Psi(t) + \eta$ with $\eta \geq 0$, then*

$$\frac{A}{N^2} \leq \frac{\eta}{2 - 3t}, \quad \frac{T}{N^3} \leq \frac{3\eta}{4}.$$

There is a nonzero error value v such that the proportion of domain points with nonzero error different from v is at most $\eta/[t(2 - 3t)]$.

Proof. The first two inequalities follow from (II.294.1). If $m_* = \max c_v$, then $M \leq sm_*$, whence $A \geq s(s - m_*)$. Divide by N^2 . $\square$

II.295 Sharp thresholds for global affine recovery

Source 67, delivered Section 4.

The local theorem makes no claim that its affine reference is nearest. This section constructs a nearby affine reference from energy alone. It also distinguishes a sharp admissible defect threshold from an exact energy envelope at every real parameter: equality with the cubic profile requires subgroup indices, which form a discrete set.

II.295.1 A scalar Fourier argument

Use the normalized Fourier transform on G , $\hat{u}(\gamma) = N^{-1} \sum_x u(x) \overline{\gamma(x)}$. For $f : G \rightarrow \mathbb{F}_3$, set $u(x) = \omega^{f(x)}$. Orthogonality in the target and in the domain gives

$$\frac{\mathcal{E}(f)}{N^3} = \frac{1 + 2 \sum_{\gamma \in \widehat{G}} |\hat{u}(\gamma)|^4}{3}, \quad \sum_{\gamma} |\hat{u}(\gamma)|^2 = 1. \quad (\text{II.295.1})$$

Lemma II.295.1 (Excluding a character of the wrong order). *If $u : G \rightarrow \{1, \omega, \omega^2\}$ and γ^3 is a nontrivial character, then*

$$|\hat{u}(\gamma)| \leq \frac{\sqrt{3}}{2}. \quad (\text{II.295.2})$$

Proof. For $z = u\bar{\gamma}$ one has $|z| = 1$ and $\mathbb{E}z^3 = \mathbb{E}\bar{\gamma}^3 = 0$. Put $r = \sqrt{3}/2$. For every real θ ,

$$r + \frac{1}{6} \cos(3\theta) - \cos \theta = \frac{2}{3} (\cos \theta - r)^2 (\cos \theta + 2r) \geq 0.$$

Apply this inequality to the phase of $e^{-i\phi}z$, choosing ϕ so that $e^{-i\phi}\mathbb{E}z$ is nonnegative real. The third moment term averages to zero, proving the bound. $\square$

For later use, if p_0, p_1, p_2 are the proportions of the three values of a ternary function, then

$$|p_0 + p_1\omega + p_2\omega^2|^2 = \frac{3(p_0^2 + p_1^2 + p_2^2) - 1}{2}. \quad (\text{II.295.3})$$

In particular, a squared modulus above $3/4$ forces one $p_j > 5/6$, using $\sum p_j^2 \leq \max p_j$. A squared modulus above $1/3$ forces one $p_j > 2/3$: with all $p_j \leq 2/3$, the maximal sum of their squares is $5/9$, attained by a permutation of $(2/3, 1/3, 0)$.

Let

$$\rho = \frac{2 - \sqrt{2}}{3}.$$

The cubic Ψ is strictly increasing on $[0, \rho]$ and $\Psi(\rho) = \Psi(1/3) = 4/9$. Also

$$\Psi(t) \geq \frac{4}{9} \quad (\rho \leq t \leq 1/3), \quad 3\rho = 2 - \sqrt{2} < \frac{2}{3}. \quad (\text{II.295.4})$$

These facts follow either by differentiation or from

$$\Psi(t) - \frac{4}{9} = \frac{2}{9} (3t - 1)(9t^2 - 12t + 2).$$

Write $D : [0, 4/9] \rightarrow [0, \rho]$ for this increasing inverse branch.

Proposition II.295.2 (Scalar recovery). *Let $f : G \rightarrow \mathbb{F}_3$.*

1. *If $\varepsilon(f) < 1/4$, there is a unique nearest affine map ℓ , and its disagreement proportion is at most $D(\varepsilon(f))$.*
2. *If G is elementary abelian of exponent three, the same conclusion holds under the weaker hypothesis $\varepsilon(f) < 4/9$.*

Proof. Put $q_\gamma = |\widehat{u}(\gamma)|^2$, $Q = \sum q_\gamma^2$, and $M = \max q_\gamma$. By (II.295.1), $Q = 1 - 3\varepsilon(f)/2$.

In the first case $Q > 5/8$. Since $Q \leq M$, one has $M > 5/8$; also $Q \leq M^2 + (1 - M)^2$. The latter expression is increasing for $M \geq 1/2$ and is $5/8$ at $M = 3/4$, so $M > 3/4$. Lemma II.295.1 shows that a character attaining M has order dividing three. It therefore equals ω^L for a homomorphism $L : G \rightarrow \mathbb{F}_3$. Equation (II.295.3), applied to $f - L$, gives an affine map $L + j$ with disagreement $t < 1/6$. Theorem II.294.1 gives $\varepsilon(f) \geq \Psi(t)$, so $t \leq D(\varepsilon(f))$. The relevant interval is increasing; moreover $\Psi(1/12) = 77/288 > 1/4$, so in fact $t < 1/12$.

In the second case $Q > 1/3$, so some $q_\gamma > 1/3$. Every character of G has order dividing three. The second consequence of (II.295.3) gives an affine map with disagreement $t < 1/3$. The local cubic and (II.295.4) force $t < \rho$, and inversion gives $t \leq D(\varepsilon(f))$.

Finally, two distinct affine maps into $\mathbb{F}_3$ agree on at most $N/3$ points: a nonconstant affine difference has equal-sized fibers, and a nonzero constant difference has no zeros. In both cases $t < \rho < 1/3$, so no two affine maps can both have disagreement at most t . In particular, a nearest affine map is unique and is the map just constructed. $\square$

II.295.2 Synchronizing all linear projections

Theorem II.295.3 (Two sharp universal recovery thresholds). *Let W be an elementary abelian 3-group and let $f : G \rightarrow W$ on a finite abelian group G . Then there is a unique nearest affine map ℓ satisfying*

$$\frac{|\{x : f(x) \neq \ell(x)\}|}{N} \leq D(\varepsilon(f)) \quad (\text{II.295.5})$$

whenever either

$$\varepsilon(f) < \frac{1}{4}, \quad \text{or} \quad G \text{ is elementary abelian of exponent three and } \varepsilon(f) < \frac{4}{9}.$$

Neither strict defect threshold can be enlarged, or replaced by a non-strict inequality, in its stated class of domains. If equality holds in (II.295.5) with positive defect, the error $f - \ell$ is precisely one nonzero constant on a subgroup coset. Such examples attain the cubic inverse at their available subgroup-index parameters.

Proof. First suppose W finite-dimensional. For each $\lambda \in W^*$, let $f_\lambda = \lambda \circ f$ and let ℓ_λ be the affine map supplied by Proposition II.295.2. A failed projected quadruple is a failed original quadruple, so $\varepsilon(f_\lambda) \leq \varepsilon(f)$ and

$$\Pr[f_\lambda \neq \ell_\lambda] \leq D(\varepsilon(f_\lambda)) \leq D(\varepsilon(f)) < \rho.$$

For λ, μ , the affine map $\ell_{\lambda+\mu} - \ell_\lambda - \ell_\mu$ is nonzero at a proportion of inputs at most $3D(\varepsilon(f)) < 3\rho < 2/3$. A nonzero affine map into $\mathbb{F}_3$ is nonzero on at least $2/3$ of the domain. Consequently

$$\ell_{\lambda+\mu} = \ell_\lambda + \ell_\mu, \quad \ell_0 = 0.$$

Finite-dimensional duality now supplies a unique map $\ell : G \rightarrow W$ with $\lambda \circ \ell = \ell_\lambda$ for every λ . Since the scalar maps are affine and linear functionals separate points, ℓ is affine.

For each nonzero vector v , exactly $2/3$ of the functionals are nonzero on v . Hence, if $t = \Pr[f \neq \ell]$,

$$\frac{2}{3}t = \mathbb{E}_\lambda \Pr[f_\lambda \neq \ell_\lambda] \leq D(\varepsilon(f)).$$

It follows that $t < (3/2)\rho < 1/3$. Apply Theorem II.294.1 and (II.295.4) once more: because $\varepsilon(f) < 4/9$, the interval $[\rho, 1/3]$ is excluded and $t \leq D(\varepsilon(f))$. Distinct affine maps into W disagree on at least $2N/3$ points, so the nearest map is unique.

If W is infinite, let V be the finite-dimensional span of $f(G)$ and apply the construction in V . Any competing affine map into W whose distance is no greater than the constructed map has distance less than $N/3$; the same affine-distance bound forces it to coincide with the constructed map. This proves global uniqueness in W , not merely uniqueness among maps into V . Equality in the inverse is equality in the local cubic and therefore has the stated coset form.

For sharpness of the first threshold, use $G = \mathbb{Z}/2\mathbb{Z}$ and $f(0) = 0, f(1) = 1$ in $\mathbb{F}_3$. There are six respected quadruples, so $\varepsilon(f) = 1/4$. Every homomorphism $\mathbb{Z}/2\mathbb{Z} \rightarrow \mathbb{F}_3$ is zero; the constants zero and one are distinct nearest affine maps, each at distance $1/2$. For the second threshold, use $G = \mathbb{F}_3$ and $f = (0, 0, 1)$. Its graph energy is 15, so its defect is $4/9$. The three affine maps interpolating each pair of its three graph points are distinct and all have distance $1/3$. Both examples embed into any nonzero elementary abelian 3-target and lift to arbitrarily large domains by a product projection. Thus the two claimed universal thresholds are sharp. $\square$

[write] 7 October 2026; sources 60–69. *Two different “4/9”* (checked at the write). Here $4/9$ bounds the *defect* $\varepsilon(f)$ on elementary abelian 3-domains; source 59’s $t < 4/9$ bounds the *error density* in the local cubic (Theorem II.250.4). The recovery threshold $1/4$ on arbitrary domains compares with source 11’s $1/9$ (Theorem F.16.2) and source 59’s $2/9$ for involution-free targets (Theorem II.251.2).

The inverse bound should be interpreted at attainable data. At positive equality parameters t is the reciprocal of a subgroup index. No assertion is made that every real ε in either interval is attained with equality.

Lemma II.295.4 (Ternary quantization on thirteen roots). *If $u : \mathbb{Z}/13\mathbb{Z} \rightarrow \{1, \omega, \omega^2\}$ and γ is a nontrivial character, then*

$$|\widehat{u}(\gamma)| \leq \frac{\sin(\pi/3)}{13 \sin(\pi/39)}, \quad |\widehat{u}(\gamma)|^2 < \frac{7}{10}. \quad (\text{II.295.6})$$

Proof. Write the thirteen values of γ in their cyclic order. For any phase ϕ , the real part of $e^{-i\phi} u(x) \overline{\gamma(x)}$ is at most the largest cosine obtained by choosing one of the three possible values of $u(x)$. Because multiplication by three permutes the residues modulo thirteen, the character phases, reduced modulo $2\pi/3$, are thirteen equally spaced points with spacing $\Delta = 2\pi/39$. Choose for each one its representative in $[-\pi/3, \pi/3]$ after the common rotation by $-\phi$. These representatives have the form $\delta + j\Delta$ for $-6 \leq j \leq 6$, with $|\delta| \leq \Delta/2$; at a boundary either representative has the same cosine. Therefore their cosine sum is

$$\sum_{j=-6}^6 \cos(\delta + j\Delta) = \cos \delta \frac{\sin(13\Delta/2)}{\sin(\Delta/2)} \leq \frac{\sin(\pi/3)}{\sin(\pi/39)}.$$

Choose ϕ to be the argument of $\widehat{u}(\gamma)$ and divide by thirteen to obtain the first bound.




Figure II.295.1: The local cubic profile and the two strict universal recovery thresholds. The profile is relative to any affine reference and is not monotone over its full positive interval. Recovery uses its increasing branch up to $\rho = (2 - \sqrt{2})/3$; the proof excludes the intervening branch before invoking the inverse. The plot is not a claim that every point on the curve is an attainable nearest-affine extremum.

Strict concavity of sine gives $13 \sin(\pi/39) > 4 \sin(\pi/12)$. Consequently

$$\left(\frac{\sin(\pi/3)}{13 \sin(\pi/39)} \right)^2 < \left(\frac{\sin(\pi/3)}{4 \sin(\pi/12)} \right)^2 = \frac{3(2 + \sqrt{3})}{16} < \frac{7}{10}.$$

The last inequality is equivalent to $\sqrt{3} < 26/15$, whose square is $675 < 676$. $\square$

Corollary II.295.5 (The sharp unrestricted-domain cutoff). *Suppose $W \neq \{0\}$ is elementary abelian of exponent three and $|G| = N \geq 13$. Every non-affine $f : G \rightarrow W$ satisfies*

$$\mathcal{E}(f) \leq N^3 - 4N^2 + 10N - 6.$$

Equality holds exactly for one-point modifications of affine maps. The universal cutoff thirteen is optimal.

Proof. For $N \geq 14$, $\Psi(1/N) \leq \Psi(1/14) < 1/4$. If a non-affine map had larger energy, Theorem II.295.3 would give affine distance less than $1/N$, impossible on an N -point set. At equality the recovered distance is at most $1/N$, and hence exactly $1/N$; the map is a one-point modification. Conversely such a modification has the displayed energy by (II.294.11).

The remaining order thirteen. For $N = 13$, the domain is $G = \mathbb{Z}/13\mathbb{Z}$ and every affine map into an elementary abelian 3-group is constant. The one-point value of the graph energy is $13^3 - 4 \cdot 13^2 + 10 \cdot 13 - 6 = 1645$.

First suppose $f : G \rightarrow \mathbb{F}_3$ has graph energy at least 1645. With Q and M as in Proposition II.295.2,

$$Q \geq \frac{1369}{2197} > \frac{29}{50}, \quad Q \leq M, \quad Q \leq M^2 + (1 - M)^2.$$

Thus $M > 7/10$: otherwise $M > 1/2$ and the last quadratic would give $Q \leq (7/10)^2 + (3/10)^2 = 29/50$. By Lemma II.295.4, the maximizing character must be trivial. Equation (II.295.3) then gives

$$\max_j \Pr[f = j] \geq \sum_j \Pr[f = j]^2 = \frac{2M + 1}{3} > \frac{4}{5}.$$

There is therefore a constant j with at most two disagreements. If there were exactly two, Theorem II.294.1 would imply

$$\varepsilon(f) \geq \Psi(2/13) = \frac{880}{2197} > \frac{552}{2197} = 1 - \frac{1645}{2197},$$

a contradiction. Thus f differs from a constant at at most one point.

For a vector-valued $f : G \rightarrow W$ with the same energy lower bound, restrict to the finite-dimensional span V of its values. Each projection $\lambda \circ f$ has at least the same graph energy and hence differs from a uniquely determined constant j_λ at at most one point. For any $\lambda, \mu \in V^*$, outside the union of their three possible error sets one has $j_{\lambda+\mu} = j_\lambda + j_\mu$. Such a point exists because $13 > 3$. Hence $\lambda \mapsto j_\lambda$ is linear, so there is $w \in V$ with $j_\lambda = \lambda(w)$ for every λ . Average the number of projected disagreements over V^* :

$$\frac{2}{3} |\{x : f(x) \neq w\}| = \mathbb{E}_\lambda |\{x : \lambda(f(x)) \neq j_\lambda\}| \leq 1.$$

The integer on the left before multiplication by $2/3$ is consequently at most one. Thus a nonconstant f with energy at least 1645 is a one-point modification of a constant. Its energy is exactly 1645 by (II.294.11). This proves both the inequality and its equality statement at $N = 13$. $\square$

[write] **7 October 2026; sources 60–69.** *Credit* (standing rule of 4 October 2026). The example at $N = 12$ that makes 13 optimal is source 59’s parity map $\mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/3\mathbb{Z}$ (Section II.252), which this source does not attribute to source 59 here, although it credits source 59’s cutoff 16 ([271]). *Checked at the write:* $(3/4) \cdot 12^3 = 1296$ and $12^3 - 4 \cdot 12^2 + 10 \cdot 12 - 6 = 1266$. This settles source 11’s cutoff for elementary abelian 3-group targets at 13 (item II.21).

The cutoff is sharp: the index-two example on $\mathbb{Z}/12\mathbb{Z}$, with values zero and one on its two parity classes, has energy 1296, whereas the one-point value is 1266. Thus no uniform cutoff at most 12 is possible for this class. Both numbers use the graph-energy normalization in the theorem.

II.295.3 Consequences for a partial frequency graph

Suppose $B \subseteq G$, $\phi : B \rightarrow W$, and its graph Γ_B has energy at least $(1 - \varepsilon)N^3$. Every completion $f : G \rightarrow W$ then has defect at most ε . Under either strict threshold in Theorem II.295.3, the completion therefore recovers an affine map with at most $D(\varepsilon)N$ disagreements on B . Moreover this affine map is independent of the chosen completion. Indeed, writing $\beta = |B|/N$, graph energy is at most $|B|^3$, so $\beta \geq (1 - \varepsilon)^{1/3}$. Two recovered affine maps would agree on at least $(\beta - 2D(\varepsilon))N$ points of B . Even at the larger threshold,

$$(5/9)^{1/3} - 2\rho > \frac{1}{3},$$

so distinct affine maps are excluded. For an exact algebraic check, $(5/9)^{1/3} > 4/5$ and $2\rho < 2/5$, giving a lower bound greater than $2/5 > 1/3$.

This is an endpoint refinement of the graph-energy output used after Gowers’s Fourier selection step. It supplies neither a mechanism for reaching this high-energy regime from small correlation nor an improvement to the complete quantitative Szemerédi iteration.

II.296 Sources 68 and 66: the energy gap, the coset range and the fourth energy level [write]

Added in the eleventh write, 7 October 2026.

The next thirteen sections are source 68’s Sections 2–4 and source 66’s Sections 2–11, on sets of nearly maximal additive energy. *Logically they continue source 45’s chapter* on energy levels and the coset range (Section II.183), source 14’s coset envelope (Theorem II.43.2) and source 57’s third level (Section II.291); they are printed here because no number printed earlier may move. Source 68’s universal inequalities, valid in every abelian group, come first; source 66’s ordered classification, for real and torsion-free sets, second. Their introductions are Sections F.73 and F.71.

What source 68 proves. For finite A of size m in any abelian group and every d , $m r_A(d)(m - r_A(d)) \leq m^3 - E(A)$, with coefficient one sharp even at positive defect, with an exact mixed-set remainder, all equality cases and a weighted form (Theorems II.297.1, II.297.7, Propositions II.297.2, II.297.8), and the multilevel bound $E(A) \leq (2m^3 + \sum_i g_i^3)/3$ from attained overlap values (Theorem II.297.4). With $\varepsilon = 1 - E(A)/m^3$: for $\varepsilon \leq 1/4$ the set $H = \{d : r_A(d) > m/2\}$ is a subgroup and a most occupied H -coset C has $|A \triangle C| \leq \tau(\varepsilon)m$, the threshold $1/4$ optimal even in cyclic groups (Theorem II.298.1, Proposition II.298.6); the closest coset is unique below the separate optimal threshold $2/9$ (Corollary II.298.4); all envelope equality cases are punctured cosets $C \setminus (z + L)$ of index at least four (Theorem II.298.5); and nested puncture repair holds with outer defect below

1/4 and inner defect at most 1/4, with all equality cases and an oriented version from one nearly saturated translation (Theorems II.299.4, II.299.5, II.299.8, Proposition II.299.11). Source 14 had the coset envelope for $\varepsilon \leq 1/100$ and source 45 for $\varepsilon < 1/9$ with the coefficient two in the pointwise gap; this answers source 45's Questions 6.6–6.8 and source 14's Question 4 in its range and equality parts (item II.31).

What source 66 proves. For a positive ordered set with Schur defect $\sigma(C) \leq r$ and $n \geq 3r + 1$, $c_i = d(i + \lambda_i)$ with (λ_i) nondecreasing, nonnegative and summing to $\sigma(C)$, so defect r has exactly $p(r)$ classes (Theorem II.302.1); $n = 2r + 1$ does not suffice (Proposition II.302.4). Reflection gives bounded-hole normal forms for endpoint extensions (Theorem II.303.2), whose energy is an exact affine function of m (Theorem II.304.1), and the sharp envelope $E(A) \leq M_m - 4r(m - 1 - r)$ with equality only for $\{0, \dots, m\} \setminus \{r\}$ (Theorem II.305.2). The four largest energies for $m \geq 11$ are M_m , $M_m - 4(m - 2)$, $M_m - 8(m - 3)$, $M_m - 8m + 20$, with one, two, three and two oriented templates (Theorems II.307.1, II.308.1, II.308.2), with a recovery window, a torsion-free transfer and an affine graph consequence (Corollary II.309.1, Theorem II.309.2, Corollary II.309.3). The first gap is implicit in the equality theory of Pollard's inequality and the first two levels are source 45's, as source 66 says.

Against source 57 (Section II.291; the two were written the same day, and neither cites the other). The partition normal form is source 57's Theorem II.292.3 with threshold $3r + 1$ in place of $5s - 1$: the same for $r = 1$ (4), and better for $r \geq 3$; for $r = 2$ source 57's optimal threshold 6 (Theorem II.292.4) is better than 7. So the optimal threshold $N(r)$ satisfies $2r + 2 \leq N(r) \leq 3r + 1$ for $r \geq 3$, with $N(1) = 4$ and $N(2) = 6$, and source 66 conjectures $N(r) = 2r + 2$ (its Conjecture V.249.1). The third level: source 66's oriented templates J_m^- , J_m^+ and C_m are source 57's two affine classes J_m and K_m (the first in its two orientations, $C_m = K_m$ symmetric); source 57's cutoff $m \geq 7$ is better than source 66's $m \geq 8$. The fourth level is new with source 66 (it answers source 57's Research question V.216.2 for the fourth level), and so are the hole energy law (its Research question V.216.3) and the endpoint envelope (item II.53).

Relation to the formal statements. No declaration concerns additive energy levels, Schur defects or coset envelopes (checked at HEAD fb9602e55); nothing in this chapter is formalized.

II.297 A sharp pointwise energy inequality

Source 68, delivered Section 2.

Throughout this section, A is a nonempty finite subset of an abelian group G , written additively. The ambient group need not be finite. All convolutions and sums in Sections II.297–II.299 use counting measure. Set

$$m = |A|, \quad r_A(d) = |A \cap (A + d)|, \quad E(A) = \sum_{d \in G} r_A(d)^2, \quad \varepsilon = \frac{m^3 - E(A)}{m^3}. \quad (\text{II.297.1})$$

The same energy is $\sum_x (\mathbf{1}_A * \mathbf{1}_A)(x)^2$, by reordering the equation $a_1 + a_2 = a_3 + a_4$. In particular, $0 \leq \varepsilon < 1$. Write $q(d) = r_A(d)/m$ and

$$\tau(u) = \frac{1 - \sqrt{1 - 4u}}{2} \quad (0 \leq u \leq 1/4).$$

Thus $\tau(u)(1 - \tau(u)) = u$ and τ is increasing.

II.297.1 An exact nonnegative remainder

Theorem II.297.1 (Universal overlap gap). *For every $d \in G$,*

$$m r_A(d)(m - r_A(d)) \leq m^3 - E(A). \quad (\text{II.297.2})$$

Equivalently, $q(d)(1 - q(d)) \leq \varepsilon$. The coefficient 1 is sharp, including at positive defect.

[write] **7 October 2026; sources 60–69.** *This answers source 45's Question V.166.7 (coefficient one, universally; source 45 had two).*

We prove a mixed version, which also identifies all losses in the proof. For finite sets A, B , define $E(A, B) = \sum_x (\mathbf{1}_A * \mathbf{1}_B)(x)^2$.

Proposition II.297.2 (Mixed-set remainder). *Suppose $a = |A| \geq b = |B| > 0$. For $d \in G$ put*

$$C = A \cap (B + d), \quad U = A \cup (B + d), \quad P = A \setminus (B + d), \quad Q = (B + d) \setminus A.$$

*Let $t = |C|$, $p = a - t$, $k = b - t$, and $w = \mathbf{1}_C * \mathbf{1}_U$, $v = \mathbf{1}_P * \mathbf{1}_Q$. Then*

$$\begin{aligned} ab^2 - E(A, B) - bt(b - t) &= pk^2 - E(P, Q) \\ &\quad + \sum_x (t - w(x))(w(x) + 2v(x)) \geq 0. \end{aligned} \quad (\text{II.297.3})$$

Consequently $E(A, B) \leq ab^2 - bt(b - t)$.

Proof. The disjoint decompositions $A = C \sqcup P$ and $B + d = C \sqcup Q$ give, using commutativity,

$$s := \mathbf{1}_A * \mathbf{1}_{B+d} = \mathbf{1}_C * \mathbf{1}_{C \sqcup P \sqcup Q} + \mathbf{1}_P * \mathbf{1}_Q = w + v.$$

Here $0 \leq w \leq t$, $0 \leq v \leq k$, and

$$\sum w = t(a + b - t), \quad \sum v = pk, \quad \sum s^2 = E(A, B).$$

Expanding the right side of (II.297.3) gives $pk^2 + t \sum w + 2t \sum v - \sum s^2$. Its constant part is $pk^2 + t^2(a + b - t) + 2tpk = ab^2 - bt(b - t)$. Finally $E(P, Q) = \sum v^2 \leq k \sum v = pk^2$, and every term of the remaining sum is nonnegative. $\square$

Taking $B = A$ proves Theorem II.297.1. In this case $p = k = m - t$ and the exact identity is

$$m^3 - E(A) - mt(m - t) = k^3 - E(P, Q) + \sum_x (t - w(x))(w(x) + 2v(x)). \quad (\text{II.297.4})$$

This union–intersection argument belongs to the classical set-transformation approach to additive inequalities. The contribution needed here is the precise energy remainder and its consequences; no claim to originate that general proof mechanism is made.

Example II.297.3 (Sharpness at positive defect). Let $L < K$ be finite subgroups with $r = [K : L] \geq 3$, let $A = K \setminus L$, and write $\ell = |L|$. Then $m = (r - 1)\ell$, and

$$r_A(d) = \begin{cases} m, & d \in L, \\ (r - 2)\ell, & d \in K \setminus L, \\ 0, & d \notin K. \end{cases}$$

Summation gives

$$\varepsilon = \frac{r-2}{(r-1)^2}. \quad (\text{II.297.5})$$

Every $d \in K \setminus L$ attains equality in (II.297.2). Translations of these examples have the same properties.

The earlier project bound was $q(d)(1-q(d)) \leq 2\varepsilon$ [280]. Thus Theorem II.297.1 answers its explicit coefficient-one question. It does so for every finite subset of every abelian group, without a small-defect assumption.

II.297.2 Several overlap levels give a stronger bound

The next result preserves more information than one application of Theorem II.297.1. It is important that its truncation function uses *sum* multiplicities and its nodes come from *difference* multiplicities.

Theorem II.297.4 (Multilevel energy bound). *Let $0 = t_0 < t_1 < \cdots < t_\ell = m$, with each internal t_i an attained value of r_A . Put $g_i = t_i - t_{i-1}$. Then*

$$E(A) \leq \frac{2m^3 + \sum_{i=1}^{\ell} g_i^3}{3}. \quad (\text{II.297.6})$$

Equivalently, for $a_i = t_i/m$,

$$\varepsilon \geq \frac{1 - \sum_{i=1}^{\ell} (a_i - a_{i-1})^3}{3}.$$

Proof. Put $z = \mathbf{1}_A * \mathbf{1}_A$ and $S(u) = \sum_x \min(z(x), u)$ for $0 \leq u \leq m$. For an attained value $t = r_A(d)$, the proof above has $0 \leq w \leq \min(s, t)$ and $\sum w = t(2m - t)$. Since s is a translate of z ,

$$S(t) \geq t(2m - t). \quad (\text{II.297.7})$$

The function S is concave, $S(0) = 0$, and $S(m) = m^2$. On each interval $[t_{i-1}, t_i]$, it is therefore at least the chord joining the corresponding values of $P(u) = 2mu - u^2$. The difference between $P(u)$ and that chord is $(u - t_{i-1})(t_i - u)$; its integral is $g_i^3/6$. Consequently

$$\int_0^m S(u) du \geq \frac{2m^3}{3} - \frac{1}{6} \sum_i g_i^3.$$

On the other hand, integrating each truncated multiplicity gives

$$\int_0^m S(u) du = \sum_x \left(mz(x) - \frac{z(x)^2}{2} \right) = m^3 - \frac{E(A)}{2}.$$

Combining the two formulas proves (II.297.6). $\square$

One internal node recovers (II.297.2), since $t^3 + (m-t)^3 = m^3 - 3mt(m-t)$. Adding a genuinely new node decreases the right side of (II.297.6): splitting a positive gap $g = u + v$ replaces g^3 by $u^3 + v^3 = g^3 - 3uv g$. Thus the full distinct overlap spectrum gives the strongest member of this family.

Example II.297.5 (The interval test). If every integer $0, 1, \dots, m$ is attained by r_A , then $E(A) \leq (2m^3 + m)/3$. The integer interval $A = \{0, \dots, m-1\} \subset \mathbb{Z}$ attains this bound, because its nonzero difference multiplicities are $m, m-1, m-1, \dots, 1, 1$. The assertion is conditional on the attained spectrum. It is not an upper bound for arbitrary sets in groups with torsion.

II.297.3 Complete equality classification for one translation

Lemma II.297.6 (The elementary three-halves lemma). *If T is finite and nonempty in an abelian group and $|T - T| < 3|T|/2$, then $T - T$ is a subgroup.*

Proof. Let $n = |T|$ and $D = T - T$. For $d = a - b \in D$, the sets $T - a$ and $T - b$ both lie in D , so $r_T(d) \geq 2n - |D| > n/2$. For $d, e \in D$, the overlap triangle inequality

$$r_T(d - e) \geq r_T(d) + r_T(e) - n \quad (\text{II.297.8})$$

therefore gives $r_T(d - e) > 0$. Thus D is closed under differences. It contains zero and is finite, so it is a subgroup. For completeness, (II.297.8) follows by intersecting $T \cap (T + d)$ and $T \cap (T + e)$ inside T and translating the resulting intersection. $\square$

Theorem II.297.7 (All positive-gap equality cases). *Suppose $0 < r_A(d) < m$. Equality in (II.297.2) holds exactly in the following cases, up to translation of A .*

1. $A = K \setminus (z + L)$, where $L < K$ are finite subgroups, $z \in K$, $[K : L] \geq 3$, and $d \in K \setminus L$.
2. $A = (a + L) \sqcup (b + L)$, where L is a finite subgroup, $b - a$ has order at least four or infinite order in G/L , and $d + L = \pm(b - a) + L$.

The second family has $\varepsilon = 1/4$. Allowing order three in that family duplicates the index-three examples in the first family. If $r_A(d) \in \{0, m\}$, equality holds precisely when A is a finite subgroup coset.

Proof. Write $t = r_A(d)$ and $k = m - t$, so $t, k > 0$. Equality in (II.297.4) forces $v(x) \in \{0, k\}$, $w(x) \in \{0, t\}$, and $v(x) > 0$ only where $w(x) = t$. Hence $s = w + v$ takes only the values $0, t, m$. The set $R = \{x : s(x) = m\}$ is nonempty, since P, Q are nonempty.

Let $L = \text{Stab}(A) = \{h : A + h = A\}$, a finite subgroup, and choose $x_0 \in R$. The equation $s(x_0) = m$ says $A = x_0 - (A + d)$. The same equation at x holds exactly when $x - x_0 \in L$; therefore $R = x_0 + L$. Both P and Q are L -periodic, and $P + Q \subset R$. Fixing one element of Q shows that P lies in a single L -coset; its nonemptiness and periodicity make it exactly that coset. The same holds for Q . Thus $|L| = k$.

Pass to G/L , and let T be the image of A , of size $n = m/k \geq 2$. After a fixed translation of its argument, the sum convolution of T has exactly one value n , and all other positive values are $n - 1$. Its total mass is n^2 , so exactly n of these latter values occur: $|T + T| = n + 1$. The full representation value also gives a reflection taking T to itself. Hence $|T - T| = |T + T| = n + 1$.

For $n \geq 3$, $n + 1 < 3n/2$, and Lemma II.297.6 makes $T - T$ a subgroup of order $n + 1$. The set T lies in one coset of this subgroup and misses one point. Lifting to G gives the first family, with index $n + 1 \geq 4$. Since $0 < t < m$, the shift belongs to $K \setminus L$. For $n = 2$, write $T = \{\bar{a}, \bar{b}\}$. Its difference has order greater than two, and the nontrivial positive overlap occurs exactly at the two signed differences. Order three gives the remaining case of the first family; all other orders give the second.

Conversely, the multiplicities in Example II.297.3 prove equality for the first family. For the second family they are $2|L|$ on L , $|L|$ on each of the two signed difference cosets, and zero elsewhere. Thus $E(A) = 6|L|^3$, $m = 2|L|$, and $\varepsilon = 1/4$.

Finally, when $t = 0$ or m , equality means $E(A) = m^3$. Counting triples shows that every $a - b + c$, for $a, b, c \in A$, then belongs to A . Translating one element to zero makes A a finite subgroup. The converse is immediate. $\square$

II.297.4 Bounded weights and minimum overlap

There is a useful extension in which the correct overlap is the minimum of two weights, rather than their product.

Proposition II.297.8 (Weighted overlap gap). *Let $f : G \rightarrow [0, 1]$ have finite support and positive mass $m = \sum_x f(x)$. Define*

$$E(f) = \sum_x (f * f)(x)^2, \quad t_f(d) = \sum_x \min(f(x), f(x - d)).$$

Then, for every d ,

$$E(f) \leq m^3 - mt_f(d)(m - t_f(d)). \quad (\text{II.297.9})$$

The multilevel bound (II.297.6) holds with arbitrary attained values of t_f and these real-valued masses. If $E(f) \geq 3m^3/4$, the finite set $\{d : t_f(d) > m/2\}$ is a subgroup. More generally, for $0 \leq f \leq M$, $M > 0$, the corresponding normalized defect is $1 - E(f)/(Mm^3)$.

Proof. Put $g(x) = f(x - d)$, $c = \min(f, g)$, $u = \max(f, g)$, $p = f - c$, and $v_0 = g - c$. Then $f = c + p$, $g = c + v_0$, $u = c + p + v_0$, and commutativity gives

$$f * g = c * u + p * v_0.$$

Let $t = \sum c$, $k = m - t$, $w = c * u$, and $v = p * v_0$. Because $u, p, v_0 \leq 1$, we have $w \leq t$, $v \leq k$. Moreover $\sum w = t(2m - t)$, $\sum v = k^2$, and $\sum (f * g)^2 = E(f)$. The expansion proving (II.297.4) therefore applies without change and proves (II.297.9). It also gives $\sum_x \min((f * f)(x), t) \geq t(2m - t)$ at each attained minimum-overlap value. The concavity and integration proof of Theorem II.297.4 uses no integrality and still applies.

Finally, $t_f(d) = m - \frac{1}{2}\|f - f(\cdot - d)\|_1$. The triangle inequality for the ℓ^1 norm gives $t_f(d - e) \geq t_f(d) + t_f(e) - m$. The two-node proof of the subgroup assertion in Theorem II.298.1 now applies verbatim. Replacing f by f/M proves the last assertion. $\square$

For indicator functions, minimum overlap equals the usual autocorrelation. For general weights it does not. In particular, the occupancy identity $\sum_{d \in H} r_A(d) = \sum_i |A \cap C_i|^2$ used below has no direct replacement with t_f ; the weighted proposition does not by itself assert the set-valued coset-distance bound.

II.298 Sharp coset rigidity and two different thresholds

Source 68, delivered Section 3.

The pointwise gap separates overlap values into two ranges. Using two attained values at once proves that the upper range is a subgroup throughout the entire real range of the quadratic inverse. This extra step is necessary: the one-value triangle argument alone only reaches defect 2/9.

II.298.1 The canonical subgroup and the optimal existence range

Theorem II.298.1 (Sharp coset existence). *Let A be as in (II.297.1), and suppose $\varepsilon \leq 1/4$. Then*

$$H = \{d \in G : r_A(d) > m/2\} \quad (\text{II.298.1})$$

is a finite subgroup. If C is any H -coset maximizing $|A \cap C|$, then

$$\delta_C := \frac{|A \triangle C|}{m} \leq \tau(\varepsilon). \quad (\text{II.298.2})$$

For $\varepsilon < 1/4$, this maximizing H -coset is unique. The threshold $1/4$ is optimal for both the subgroup assertion and the guarantee that some subgroup coset has distance at most $m/2$. Optimality already holds among finite cyclic ambient groups.

Proof of the subgroup assertion. The set H contains zero, is symmetric, and is finite, because $H \subset A - A$. Suppose $d, e \in H$ but $q(d - e) \leq 1/2$. Set $c = q(d - e)$ and $b = \min(q(d), q(e))$. The overlap triangle inequality gives

$$0 < c \leq \frac{1}{2} < b \leq \frac{1 + c}{2}.$$

Both c and b are attained normalized multiplicities. Theorem II.297.4, with these two internal nodes, gives

$$\varepsilon \geq F(c, b) := \frac{1 - c^3 - (b - c)^3 - (1 - b)^3}{3}.$$

For fixed c , $\partial F / \partial b = (1 - c)(1 + c - 2b)$. It is positive before the upper endpoint of the displayed interval. As $b > 1/2$ and $c > 0$,

$$F(c, b) > F(c, 1/2) = \frac{1}{4} + \frac{c(1/2 - c)}{2} \geq \frac{1}{4},$$

a contradiction, including when $\varepsilon = 1/4$. Thus H is closed under differences and is a subgroup. $\square$

Proof of the distance bound. Write $a = \tau(\varepsilon)$, $\beta = 1 - a$, and $h = |H|$. Theorem II.297.1 implies

$$q(d) \geq \beta \quad (d \in H), \quad q(d) \leq a \quad (d \notin H).$$

These inequalities remain valid at $a = \beta = 1/2$. Let $p_i = |A \cap C_i|/m$ be the occupancies of the H -cosets, and put $p_* = \max_i p_i$, $S_2 = \sum_i p_i^2$. Then $\sum_i p_i = 1$, $S_2 \leq p_*$, and

$$\sum_{d \in H} q(d) = mS_2, \quad \sum_{d \in G} q(d) = m, \quad \varepsilon = \frac{1}{m} \sum_{d \in G} q(d)(1 - q(d)).$$

For $d \in H$, $q(1 - q) \geq \beta(1 - q)$. For $d \notin H$, $q(1 - q) \geq \beta q$. Summing these bounds gives

$$\varepsilon \geq \beta \left(\frac{h}{m} + 1 - 2S_2 \right). \quad (\text{II.298.3})$$

For a maximizing coset, $\delta_C = h/m + 1 - 2p_*$. Since $\varepsilon = a\beta$,

$$\delta_C + 2(p_* - S_2) \leq a. \quad (\text{II.298.4})$$

This proves (II.298.2).

To prove uniqueness among H -cosets when $a < 1/2$, bound $q^2 \leq q$ on H and $q^2 \leq aq$ off H . Then

$$1 - \varepsilon = \frac{1}{m} \sum_d q(d)^2 \leq S_2 + a(1 - S_2), \quad S_2 \geq 1 - \frac{\varepsilon}{1 - a} = 1 - a > \frac{1}{2}.$$

Thus $p_* \geq S_2 > 1/2$, and only one H -coset can attain p_* . The optimality assertion is proved in Proposition II.298.6 below. $\square$

[write] **7 October 2026; sources 60–69.** Source 14's envelope held for $\varepsilon \leq 1/100$ (Theorem II.43.2) and source 45's for $\varepsilon < 1/9$ (Theorem II.185.4); $\tau(1/4) = 1/2$ and $\tau(2/9) = 1/3$ (checked at the write).

The same calculation yields an exact identity that will be useful for stability and equality:

$$a - \delta_C = 2(p_* - S_2) + \frac{1}{m\beta} \left(\sum_{d \in H} (q(d) - \beta)(1 - q(d)) + \sum_{d \notin H} q(d)(a - q(d)) \right). \quad (\text{II.298.5})$$

Every summand is nonnegative. This records both the distribution of mass between cosets and the deviation of overlaps from the two extreme levels.

Corollary II.298.2 (A sharper outlier budget). *Assume $\varepsilon < 1/4$, let C be the canonical coset from Theorem II.298.1, and put $x = |A \setminus C|/m$, $\delta = \delta_C$. Then*

$$\delta + 2x(1 - 2x) \leq \tau(\varepsilon) < \frac{1}{2}, \quad x < \frac{1}{4}, \quad \delta < \frac{1}{2}. \quad (\text{II.298.6})$$

Proof. The largest occupancy is $1 - x > 1/2$, and $S_2 \leq (1 - x)^2 + x^2$. Substitute in (II.298.4). Since $\delta \geq x$, this gives $3x - 4x^2 < 1/2$. On $1/4 \leq x < 1/2$ the left side is at least $1/2$, so $x < 1/4$. $\square$

II.298.2 The sharp uniqueness range

Here and below a subgroup coset means a coset of a *finite* subgroup, even if the ambient group is infinite. Define

$$\delta(A) = \min_C \frac{|A \triangle C|}{m}.$$

The minimum exists: the set of possible unnormalized distances is a nonempty set of nonnegative integers.

Lemma II.298.3 (Separation of distinct cosets). *Distinct subgroup cosets C, C' satisfy*

$$|C \triangle C'| \geq \min(|C|, |C'|).$$

Proof. If their intersection is empty the claim is immediate. Otherwise their intersection is a coset of the intersection of their underlying subgroups. If one coset properly contains the other, the subgroup index is at least two, and the difference has size at least that of the smaller coset. If neither contains the other, the intersection has size at most half of each; hence the symmetric difference is at least the larger size. $\square$

Corollary II.298.4 (Optimal uniqueness threshold). *Any coset C with $|A \triangle C| \leq m/3$ is a closest coset. If $|A \triangle C| < m/3$, it is the unique closest coset. In particular, if $\varepsilon < 2/9$, the canonical coset is the unique closest coset and*

$$\delta(A) = \delta_C \leq \tau(\varepsilon) < 1/3. \quad (\text{II.298.7})$$

The defect threshold $2/9$ is sharp for uniqueness.

Proof. Suppose both C and C' have distance at most ρm from A . Then $|C|, |C'| \geq (1 - \rho)m$, whereas $|C \triangle C'| \leq 2\rho m$. If they are distinct, Lemma II.298.3 implies $\rho \geq 1/3$. If one distance is $m/3$ and the other is strictly smaller, the upper bound is strictly below $2m/3$ while both sizes are

at least $2m/3$, again impossible. This proves both distance assertions. Now use $\tau(2/9) = 1/3$ and Theorem II.298.1.

For sharpness, take $A = \{0, 1, 2\} \subset \mathbb{Z}/4\mathbb{Z}$. Its difference multiplicities are 3, 2, 2, 2, so $E(A) = 21$ and $\varepsilon = 2/9$. Both $\mathbb{Z}/4\mathbb{Z}$ and $\{0, 2\}$ are closest cosets, each at distance $1 = m/3$. $\square$

Uniqueness among cosets of the *canonical* subgroup persists to $\varepsilon < 1/4$. Uniqueness when the underlying subgroup is also allowed to vary stops at $2/9$. These are different assertions, and both thresholds have a role in the subsequent repair theorem.

II.298.3 All canonical-envelope equality cases

Theorem II.298.5 (Equality in the coset envelope). *Use the canonical coset convention of Theorem II.298.1.*

1. If $\varepsilon = 0$, then $A = C$ is a subgroup coset.
2. For $0 < \varepsilon < 1/4$, equality $\delta_C = \tau(\varepsilon)$ holds exactly when

$$A = C \setminus (z + L), \quad L < H, \quad z + L \subset C, \quad r = [H : L] \geq 4.$$

Here $\delta_C = 1/(r - 1)$ and $\varepsilon = (r - 2)/(r - 1)^2$. These examples also satisfy $\delta(A) = \delta_C$.

3. For $\varepsilon = 1/4$, equality $\delta_C = 1/2$ holds exactly when A is the union of two cosets of $H = \text{Stab}(A)$ whose difference has order greater than two in G/H . There are exactly two maximizing H -cosets.

In the uniqueness range $\varepsilon < 2/9$, the positive equality cases in part 2 have $r \geq 5$.

Proof. The case $\varepsilon = 0$ follows from the maximal-energy characterization in Theorem II.297.7. Suppose $0 < \varepsilon < 1/4$ and equality holds. Equation (II.298.5) forces $p_* = S_2$. Since $p_* - S_2 = \sum_i p_i(p_* - p_i)$, every nonzero occupancy equals p_* . Also $S_2 \geq \beta > 1/2$, so there is only one occupied coset: $A \subset C$. The other remainder terms show that $q(d) \in \{\beta, 1\}$ for $d \in H$.

Let $L = \text{Stab}(A) = \{d : q(d) = 1\}$ and write $\ell = |L|$. Since $\delta_C = a$ and $A \subset C$, $h/m = 1 + a$. The identity $\sum_{d \in H} q(d) = m$ now becomes

$$\ell + (h - \ell)\beta = m, \quad \ell a = m - \beta h = ma^2.$$

Therefore $\ell = ma = |C \setminus A|$. The missing set is L -periodic and has size $|L|$, so it is a single L -coset. Moreover $h/\ell = (1 + a)/a = 1 + 1/a$ is an integer greater than three: $r \geq 4$. Conversely Example II.297.3 gives the stated parameters. For $r \geq 4$ its overlaps outside L but inside H are strictly above $m/2$, so the ambient coset C is indeed the canonical one. Its distance is at most $m/3$, and Corollary II.298.4 proves it is a closest coset.

At $\varepsilon = 1/4$, the same occupancy argument gives one or two nonzero occupancies, all equal. If there were one, the zero remainder would force $q(d) = 1$ for all $d \in H$, since $q(d) > 1/2$ there. Then $h = m$ and $\delta_C = 0$, a contradiction. There are therefore two occupancies, each $1/2$. As $\delta_C = h/m$, equality makes $h = m/2$; each occupied coset is full. Their quotient difference cannot have order two, since their union would then itself be a subgroup coset of maximal energy. For order greater than two, the multiplicities are m on H and $m/2$ on the two signed difference cosets. They give $\varepsilon = 1/4$, $\delta_C = 1/2$, and $H = \text{Stab}(A)$, as required. $\square$

The curve $\tau(\varepsilon)$ is attained at the indicated discrete parameters. In particular, there is no canonical-envelope equality for $2/9 < \varepsilon < 1/4$. This arithmetic gap is relevant to any proposed continuous refinement of the bound.

II.298.4 Why the endpoint cannot be increased

Proposition II.298.6 (Sharpness of the quarter-defect threshold). *For every $\varepsilon_0 > 1/4$ there is a finite subset of a finite cyclic group with defect between $1/4$ and ε_0 for which $\{d : r_A(d) > |A|/2\}$ is not a subgroup and $\delta(A) > 1/2$.*

Proof. First let K be a finite abelian group of order $h \geq 2$, and work in $K \times \mathbb{Z}$ with

$$A = (K \times \{0, 1\}) \cup \{(0, 2)\}, \quad m = 2h + 1.$$

At second coordinate zero the difference multiplicity is m at zero and $2h$ at the other $h - 1$ points. At each of the two levels ± 1 , every first coordinate has multiplicity $h + 1$. At each level ± 2 , every first coordinate has multiplicity 1. Hence

$$\begin{aligned} E(A) &= (2h + 1)^2 + (h - 1)(2h)^2 + 2h(h + 1)^2 + 2h \\ &= 6h^3 + 4h^2 + 8h + 1, \\ \varepsilon - \frac{1}{4} &= \frac{20h^2 - 14h - 1}{4(2h + 1)^3} > 0. \end{aligned}$$

The last expression tends to zero as $h \rightarrow \infty$. For $d = (0, 1)$, $r_A(d) = h + 1 > m/2$ but $r_A(2d) = 1 \leq m/2$. Every finite subgroup coset in $K \times \mathbb{Z}$ lies at one integer level, so its maximum possible improvement over the empty-set distance is h . The nearest such cosets are the two full levels, and

$$\delta(A) = \frac{h + 1}{2h + 1} > \frac{1}{2}.$$

To obtain a finite cyclic example, take $K = \mathbb{Z}/h\mathbb{Z}$ and a prime $p > 4h + 3$, and replace the integer coordinate by $\mathbb{Z}/p\mathbb{Z}$. The five difference levels remain distinct, so the energy is unchanged. A subgroup with nonzero projection onto $\mathbb{Z}/p\mathbb{Z}$ has size at least p , and each of its cosets has distance at least $p - m > h + 1$ from A . A subgroup with zero projection lies in one level, where the preceding minimum remains $h + 1$. Finally p is coprime to h , so $\mathbb{Z}/h\mathbb{Z} \times \mathbb{Z}/p\mathbb{Z}$ is cyclic. Choosing h sufficiently large proves the assertion. $\square$

II.299 Sharp puncture repair and a single near-saturated translation

Source 68, delivered Section 4.

The coset theorem determines an ambient structure. The next question is whether a set close to its quadratic error envelope is close to a coset with one subcoset removed. Source 45 of the repository report proves the profile below when both the outer energy defect and the normalized inner excess are less than $1/9$ [280]. We retain that profile with both defects less than $1/4$. We also show that near equality in the energy inequality at one translation specifies which subgroup the repaired hole must avoid.

The distinction between a canonical coset and a closest coset matters in this section. Unless a smaller range is specified, the outer coset is the one constructed by Theorem II.298.1. Above defect 2/9 it is not silently identified with a closest coset. All minimizations over holes keep this outer coset fixed.

II.299.1 A stronger remainder around a proposed coset

Retain $m, E(A), \varepsilon, q$ from (II.297.1). For any finite subgroup coset $C = c + H$, put

$$A_0 = A \cap C, \quad T = A \setminus C, \quad K = C \setminus A, \quad b = |A_0|, \quad t = |T|, \quad u = |K|, \quad (\text{II.299.1})$$

and define

$$x = \frac{t}{m}, \quad \delta = \frac{t+u}{m}.$$

An empty auxiliary set has energy zero.

Proposition II.299.1 (Outlier and missing-set remainder). *For every proposed coset C ,*

$$E(A) \leq E(A_0) + 6bt^2 + t^3, \quad (\text{II.299.2})$$

$$E(A_0) = b^3 - bu(b-u) - (u^3 - E(K)). \quad (\text{II.299.3})$$

Consequently,

$$\varepsilon \geq \delta(1-\delta) + xB(\delta, x) + \frac{u^3 - E(K)}{m^3}, \quad B(\delta, x) = 2 + \delta^2 - (8 + \delta)x + 6x^2. \quad (\text{II.299.4})$$

The polynomial B need not be positive for an arbitrary proposed coset.

Proof. Count the ordered triples $(z_1, z_2, z_3) \in A^3$ whose alternating sum $z_1 - z_2 + z_3$ lies in A , according to the number of their entries in T . Triples with no such entry contribute $E(A_0)$. If exactly one entry lies in T , the output lies outside C , hence in T . There are three choices of its position. Fixing the input in T , the output in T , and one of the two inputs in A_0 determines the remaining input. These cases contribute at most $3bt^2$. Triples with exactly two entries in T contribute at most all $3bt^2$ possible input triples, and those with three contribute at most t^3 . This proves (II.299.2).

Translate C to H and write $h = |H| = b + u$. Expanding the four indicator factors in the additive-quadruple count, using $\mathbf{1}_{A_0} = \mathbf{1}_H - \mathbf{1}_K$, gives the exact identity

$$E(A_0) = h^3 - 4h^2u + 6hu^2 - 4u^3 + E(K).$$

For example, a term with two prescribed entries in K has hu^2 solutions: one of the remaining entries is free in H , and the last is determined there. Substitution of $h = b + u$ gives (II.299.3). Finally use $b/m = 1 - x$ and $u/m = \delta - x$ in the first two assertions and expand. The resulting polynomial is (II.299.4). $\square$

Remark II.299.2 (Comparison with the predecessor remainder). Source 45 bounds the outlier contribution by $6bt^2 + 5t^3$; its resulting polynomial has $2x^2$ in place of $6x^2$. Counting input triples gives the stronger form above. This counting argument is already implicit in Source 14's proof of the quadratic coset envelope. The refinement used here combines that argument with Source 45's exact missing-set identity; we credit those predecessor ingredients [278, 280].

II.299.2 The canonical outlier budget

Suppose now that $\varepsilon < 1/4$. Let

$$a = \tau(\varepsilon), \quad H = \{d : q(d) > 1/2\},$$

and let C be the H -coset with largest intersection with A , as in Theorem II.298.1. Define x, δ by (II.299.1), and put

$$\eta = \varepsilon - \delta(1 - \delta). \quad (\text{II.299.5})$$

Here δ is the distance to this specified C .

Lemma II.299.3 (Uniform canonical budget). *One has*

$$0 \leq x < \frac{1}{4}, \quad x \leq \delta \leq a < \frac{1}{2}, \quad \eta \geq \frac{3}{8}x + \frac{u^3 - E(K)}{m^3}. \quad (\text{II.299.6})$$

Proof. Let p_j be the proportions of A in the H -cosets, and write $p_* = 1 - x$ for their maximum. The canonical-coset argument gives

$$\sum_j p_j^2 \geq 1 - a > \frac{1}{2}, \quad p_* \geq \sum_j p_j^2,$$

so $x < 1/2$ and the maximizing coset is unique. Keeping the occupancy loss in the same argument gives

$$\delta + 2 \left(p_* - \sum_j p_j^2 \right) \leq a. \quad (\text{II.299.7})$$

For completeness, the pointwise gap gives $q \geq 1 - a$ on H and $q \leq a$ off H , whence

$$\varepsilon = \frac{1}{m} \sum_d q(d)(1 - q(d)) \geq (1 - a) \left(\frac{|H|}{m} + 1 - 2 \sum_j p_j^2 \right).$$

Use $\varepsilon = a(1 - a)$ and $\delta = |H|/m + 1 - 2p_*$ to obtain (II.299.7).

Since $\sum_j p_j^2 \leq (1 - x)^2 + x^2$, this implies

$$\delta + 2x(1 - 2x) \leq a < \frac{1}{2}.$$

Also $\delta \geq x$. Thus $3x - 4x^2 < 1/2$. This is impossible for $1/4 \leq x < 1/2$, proving $x < 1/4$. For $\delta \geq x$,

$$\frac{\partial B}{\partial \delta} = 2\delta - x \geq 0, \quad B(\delta, x) \geq 2 - 8x + 6x^2 > \frac{3}{8}.$$

The last polynomial decreases on $[0, 1/4]$ and has value $3/8$ at $1/4$. Insert this bound into (II.299.4). $\square$

II.299.3 The exact puncture profile below the quarter threshold

For the fixed canonical C , define

$$\Delta_C(A) = \frac{1}{m} \min_{\substack{D \subseteq C \text{ a nonempty} \\ \text{subgroup coset}}} |A \triangle (C \setminus D)|. \quad (\text{II.299.8})$$

The theorem below proves that an admissible hole exists under its hypotheses.

Theorem II.299.4 (Sharp canonical puncture profile). *Suppose $\varepsilon < 1/4$, let C be its canonical coset, and assume*

$$\delta = \frac{|A \triangle C|}{m} > 0, \quad 0 \leq s = \frac{\eta}{\delta^3} < \frac{1}{4}, \quad \eta = \varepsilon - \delta(1 - \delta). \quad (\text{II.299.9})$$

Then $K = C \setminus A$ is nonempty, a closest subgroup coset to K exists, and every such closest coset is a nonempty proper subcoset of C . Moreover,

$$\Delta_C(A) \leq \delta\tau(s). \quad (\text{II.299.10})$$

If $s < 2/9$, the minimizing hole is unique. Separately, if $\varepsilon < 2/9$, C is the unique closest outer coset to A .

Proof. Put $c = 3/8$. The bound $\eta \geq cx$ holds even if K is initially empty. Since

$$x \leq \frac{\eta}{c} = \frac{8}{3}s\delta^3 < \frac{2}{3}\delta^3 < \frac{\delta}{6},$$

we have $x < \delta$, so K is nonempty. Set

$$k = \frac{|K|}{m} = \delta - x, \quad \zeta = 1 - \frac{E(K)}{|K|^3}.$$

Lemma II.299.3 gives the coupled inequality

$$\eta \geq cx + k^3\zeta. \quad (\text{II.299.11})$$

For $0 \leq y \leq \eta/c$ define

$$z(y) = \frac{\eta - cy}{(\delta - y)^3}.$$

Its denominator is positive, and

$$z'(y) = \frac{-c\delta + 3\eta - 2cy}{(\delta - y)^4} < 0. \quad (\text{II.299.12})$$

Indeed $3\eta = 3s\delta^3 < (3/4)\delta^3 < (3/16)\delta < c\delta$. It follows that

$$0 \leq \zeta \leq z(x) \leq z(0) = s < \frac{1}{4}.$$

Apply Theorem II.298.1 to K within C , after translating C to H . It supplies a subgroup coset $D \subset C$ such that

$$|K \triangle D| \leq \tau(\zeta)|K|. \quad (\text{II.299.13})$$

This D is proper. Writing $h = |C|$, we have

$$\frac{|K|}{h} = \frac{\delta - x}{1 + \delta - 2x} < \frac{1}{3},$$

because $2\delta - x < 1$. Since $\tau(\zeta) < 1/2$, $|D| < (3/2)|K| < h/2$.

For any hole inside C , the outliers T contribute exactly t edits. Hence

$$\Delta_C(A) \leq x + k\tau(\zeta) \leq x + (\delta - x)\tau(z(x)). \quad (\text{II.299.14})$$

To optimize this coupled bound without losing a leading constant, set

$$F(y) = y + (\delta - y)\tau(z(y)), \quad b_y = \tau(z(y)).$$

Using $b_y(1 - b_y) = z(y)$, direct differentiation gives

$$F'(y) = \frac{1 - b_y^2 - c/(\delta - y)^2}{1 - 2b_y} < 0. \quad (\text{II.299.15})$$

The denominator is positive, while $\delta - y \leq \delta < 1/2$ gives $c/(\delta - y)^2 > 3/2$. Thus

$$\Delta_C(A) \leq F(x) \leq F(0) = \delta\tau(s),$$

as claimed.

We justify the assertion about closest holes. A closest subgroup coset to K exists because the nonnegative integer distances have a minimum. It must intersect K : a disjoint coset has distance at least $|K| + 1$, whereas a singleton in K has distance $|K| - 1$. If a closest coset E were not contained in C , its nonempty intersection $E \cap C$ would be a subgroup coset and would strictly decrease its distance to K . Finally $E \neq C$, since $|C| - |K| > 2|K|$ and a singleton is closer. Therefore all closest cosets are admissible holes, and they minimize (II.299.8).

If $s < 2/9$, then $\zeta \leq s < 2/9$ and Corollary II.298.4 applied to K makes its closest coset unique. The assertion about the outer coset follows from that corollary applied to A . If $s = 0$, the budget gives $x = 0$ and $E(K) = |K|^3$, so K itself is a coset and the repair is exact. $\square$

Theorem II.299.5 (All equality cases in the puncture profile). *Under the hypotheses of Theorem II.299.4, if $s > 0$, equality in (II.299.10) holds precisely when*

$$A = (C \setminus D) \cup J, \quad (\text{II.299.16})$$

where $J \subsetneq D \subsetneq C$ are nonempty subgroup cosets and the subgroup of J has index $r \geq 4$ in the subgroup of D . The assertion retains the specified canonical outer coset C . For $s = 0$, one has $A = C \setminus D$ exactly.

Proof. Suppose first that $s > 0$ and equality holds. In the proof of Theorem II.299.4, use the canonical inner coset supplied by Theorem II.298.1. The strict decrease in (II.299.15) forces $x = 0$. Thus $A \subset C$, $|K|/m = \delta$, and the exact complement identity gives

$$\eta = \delta^3 \zeta, \quad \zeta = s.$$

The chain of inequalities can now be an equality only if the canonical inner coset realizes equality in its own quadratic envelope. Theorem II.298.5 says exactly that $K = D \setminus J$, where the subgroup index is $r \geq 4$.

Conversely, suppose the displayed configuration holds under the theorem's hypotheses. Write $\ell = |J|$, so $|K| = (r-1)\ell$. The inner energy identity gives

$$\zeta = \frac{r-2}{(r-1)^2}, \quad \frac{|K \triangle D|}{|K|} = \frac{1}{r-1} = \tau(\zeta) \leq \frac{1}{3}.$$

Corollary II.298.4, including its endpoint minimality assertion, shows that D is a closest coset to K ; it is unique when $r \geq 5$. The complement identity then yields $\zeta = s$ and

$$\Delta_C(A) = \frac{\ell}{m} = \delta\tau(s).$$

The $s = 0$ case was proved at the end of Theorem II.299.4. □

Remark II.299.6 (What is sharp). The exact profile is attained at the discrete parameters supplied by the subgroup indices. This does not assert attainment at every real s . For $s > 0$, it can be written

$$\delta\tau(s) = \frac{\tau(s)}{s} \frac{\eta}{\delta^2}, \quad \frac{\tau(s)}{s} = \frac{1}{1-\tau(s)} = 1 + s + 2s^2 + O(s^3). \quad (\text{II.299.17})$$

The examples below prove that the uniform asymptotic leading constant one cannot be reduced.

II.299.4 A translation gap determines the orientation of the hole

For every $d \in G$ define

$$\theta_d = \varepsilon - q(d)(1 - q(d)) \geq 0. \quad (\text{II.299.18})$$

The nonnegativity is Theorem II.297.1. The excess η measures closeness to the coset envelope; θ_d measures closeness to equality at one specified translation.

Proposition II.299.7 (Exact translation-gap decomposition). *Suppose $\varepsilon < 1/4$, and use its canonical $C = c + H$. For $d \in H$ put*

$$z_d = \frac{r_K(d) + r_T(d)}{m}.$$

Then

$$q(d) = 1 - \delta + z_d, \quad \theta_d = \eta + (1 - 2\delta)z_d + z_d^2. \quad (\text{II.299.19})$$

If $d \notin H$ and $\delta > 0$, then

$$\theta_d \geq \frac{\delta}{8}. \quad (\text{II.299.20})$$

The constant in this last separation estimate is not asserted optimal.

Proof. A difference in H cannot connect C to another H -coset. Therefore

$$r_A(d) = r_{A_0}(d) + r_T(d), \quad r_{A_0}(d) = |H| - 2u + r_K(d).$$

Since $|H| - 2u = m - t - u = m(1 - \delta)$, this proves the first identity in (II.299.19). Substitute it into (II.299.18) to obtain the second.

If $d \notin H$, every pair counted by $r_A(d)$ has at least one endpoint in T , so $q(d) \leq 2x < 1/2$. The function $v(1-v)$ increases on $[0, 1/2]$. Thus (II.299.4) implies

$$\begin{aligned} \theta_d &\geq R(\delta, x), \\ R(\delta, x) &= \delta - \delta^2 + \delta^2 x - (4 + \delta)x^2 + 6x^3. \end{aligned} \quad (\text{II.299.21})$$

We claim $R(\delta, x) \geq \delta/8$ on $0 \leq x \leq \delta \leq 1/2$, $x \leq 1/4$. For fixed x , $R(\delta, x) - \delta/8$ is concave in δ , so its minimum on $[x, 1/2]$ occurs at an endpoint. At the first,

$$R(x, x) = x(1 - 2x)(1 - 3x) \geq \frac{x}{8}.$$

At the second,

$$R(1/2, x) - \frac{1}{8} = \frac{(1 - 4x)(1 + 6x - 12x^2)}{8} \geq 0,$$

which is stronger than the required $R(1/2, x) \geq 1/16$. This proves (II.299.20). $\square$

For a subgroup coset D , write L_D for its underlying subgroup. For the fixed outer coset C , define the oriented repair distance

$$\Delta_{C,d}(A) = \frac{1}{m} \min_{\substack{D \subseteq C \text{ a nonempty subgroup coset} \\ d \notin L_D}} |A \triangle (C \setminus D)|. \quad (\text{II.299.22})$$

The following theorem proves that the minimizing class is nonempty.

Theorem II.299.8 (Sharp repair from one nearly saturated translation). *Suppose $\varepsilon < 1/4$, let $C = c + H$ be the canonical coset, and assume*

$$\delta = \frac{|A \triangle C|}{m} > 0, \quad 0 \leq \sigma = \frac{\theta_d}{\delta^3} < \frac{1}{4}. \quad (\text{II.299.23})$$

Then $d \in H$. Every closest subgroup coset D to $K = C \setminus A$ is a proper subcoset of C and satisfies $d \notin L_D$. In particular,

$$\Delta_{C,d}(A) = \Delta_C(A) \leq \delta \tau \left(\frac{\theta_d}{\delta^3} \right). \quad (\text{II.299.24})$$

If $\sigma < 2/9$, the minimizing hole is unique.

For $\sigma > 0$, equality in (II.299.24) holds precisely for the nested configurations in (II.299.16), with inner index $r \geq 4$ and $d \notin L_D$. For $\sigma = 0$, the repair is exact: $A = C \setminus D$ and $d \in H \setminus L_D$.

Proof. If $d \notin H$, Proposition II.299.7 gives $\theta_d \geq \delta/8$. But the hypotheses imply

$$\theta_d < \frac{\delta^3}{4} < \frac{\delta}{16},$$

a contradiction. Hence $d \in H$, and (II.299.19) gives $\eta \leq \theta_d$. Theorem II.299.4 applies with $s = \eta/\delta^3 \leq \sigma$. For any minimizing hole D , put $e = \Delta_C(A)$. Then

$$e \leq \delta \tau(\sigma) < \frac{\delta}{2}. \quad (\text{II.299.25})$$

Suppose $d \in L_D$. Then $D + d = D$, and intersecting $K \cap D$ with its translate within D gives

$$r_K(d) \geq 2|K \cap D| - |D| = |K| - |K \triangle D|.$$

Since $e = x + |K \triangle D|/m$, division by m yields $z_d \geq r_K(d)/m \geq \delta - e > \delta/2$. The exact gap identity therefore implies

$$\theta_d \geq z_d^2 > \frac{\delta^2}{4},$$

contradicting $\theta_d < \delta^3/4 < \delta^2/8$. Thus every minimizing hole avoids d in its subgroup. All are admissible in (II.299.22), proving the equality of the two minima and the bound. If $\sigma < 2/9$, then $s < 2/9$ and the uniqueness assertion of Theorem II.299.4 applies.

Suppose next that $\sigma > 0$ and the profile is attained. The chain

$$\Delta_C(A) \leq \delta\tau(\eta/\delta^3) \leq \delta\tau(\theta_d/\delta^3)$$

forces $\eta = \theta_d$ and equality in the puncture profile, by strict monotonicity of τ . Theorem II.299.5 gives $A = (C \setminus D) \cup J$ with inner index $r \geq 4$. In this representation $r_K(d) > 0$ whenever $d \in L_D$: indeed $r_K(d) \geq |D| - 2|J| = (r-2)|J| > 0$. The equality $\theta_d = \eta$ and (II.299.19) therefore force $d \notin L_D$.

Conversely, in such a configuration with $d \notin L_D$, the hole $K = D \setminus J$ is disjoint from $K + d$ and there are no outliers. Thus $z_d = 0$, $\theta_d = \eta$, and Theorem II.299.5 gives equality in (II.299.24). If $\sigma = 0$, the same identity gives $\eta = 0$ and the exact punctured-coset representation, with the orientation already established. $\square$

II.299.5 Attainment, a nonunique hole, and the sharp leading constant

Example II.299.9 (A sharp model in the enlarged range). Let $G = \mathbb{Z}/12\mathbb{Z}$, $D = \{0, 3, 6, 9\}$, $J = \{0\}$, and $A = (G \setminus D) \cup J$. The canonical outer coset is $C = G$. Here

$$m = 9, \quad K = \{3, 6, 9\}, \quad \delta = \frac{1}{3}, \quad \eta = \frac{2}{243}, \quad \varepsilon = \frac{56}{243} \in \left(\frac{2}{9}, \frac{1}{4}\right).$$

Also $s = \eta/\delta^3 = 2/9$, and

$$\Delta_C(A) = \frac{1}{9} = \delta\tau(2/9).$$

The four-point hole D and the two-point coset $\{3, 9\}$ are both closest to K , each at one edit. Thus uniqueness of the hole would be false in the larger range. With $d = 1$, both holes avoid d in their subgroup, $\theta_d = \eta$, and the oriented profile is attained too.

Proposition II.299.10 (Exact attainment and optimal leading constant). *Fix an integer $R \geq 10$. For every integer $r \geq 4$, let $G = \mathbb{Z}/(Rr)\mathbb{Z}$, let D be its subgroup of order r , and set*

$$J = \{0\}, \quad A = (G \setminus D) \cup J, \quad d = 1.$$

Then all hypotheses of Theorem II.299.8 hold, its profile is attained, and

$$\begin{aligned} m &= (R-1)r + 1, & \delta &= \frac{r-1}{m}, \\ \theta_d = \eta &= \frac{(r-1)(r-2)}{m^3}, & \sigma &= \frac{r-2}{(r-1)^2}, \\ \Delta_{G,d}(A) &= \frac{1}{m} = \delta\tau(\sigma). \end{aligned} \tag{II.299.26}$$

Consequently the best uniform asymptotic leading constant in a bound of the form

$$\Delta_{C,d}(A) \leq (C_0 + o(1)) \frac{\theta_d}{\delta^2} \quad \text{as } \theta_d/\delta^3 \rightarrow 0$$

is $C_0 = 1$. The corresponding statement with η and Δ_C also has leading constant one.

Proof. The missing set is $K = D \setminus \{0\}$. Its maximal-coset complement calculation gives the parameters in (II.299.26). Since $d = 1 \notin D$, $r_K(d) = 0$ and hence $\theta_d = \eta$. Moreover $\delta < 1/(R-1) \leq 1/9$ and $\eta < \delta^2$, so

$$\varepsilon = \delta(1 - \delta) + \eta < \delta < \frac{1}{9}.$$

The same intersection formula shows $q(h) \geq 1 - \delta > 1/2$ for every $h \in G$, so the canonical outer coset is indeed G . For $r \geq 4$ we have $\sigma \leq 2/9 < 1/4$ and $\tau(\sigma) = 1/(r-1)$. The equality theorems give the one-edit repair. Finally,

$$\frac{\Delta_{G,d}(A)}{\theta_d/\delta^2} = \frac{r-1}{r-2} \longrightarrow 1, \quad \sigma \longrightarrow 0,$$

proving optimality of the leading constant. □

The new translation statement supplies information not present in an unoriented coset approximation: a nearly saturated overlap identifies a hole whose subgroup excludes the tested direction. Its proof retains the actual gap at that direction, the missing-set energy, and the outlier mass. It remains a local structural theorem; no improved global Szemerédi bound is inferred from it alone.

II.299.6 The inner endpoint and the strict outer threshold

The inner quarter threshold can be included. It has additional equality cases, so we record it separately from the strict-range theorems.

Proposition II.299.11 (Endpoint extension and its equality cases). *Keep the strict outer hypothesis $\varepsilon < 1/4$ and the canonical coset convention. The puncture conclusion (II.299.10) remains valid for*

$$0 \leq s = \eta/\delta^3 \leq 1/4,$$

and the oriented conclusion (II.299.24), including admissibility of every minimizing hole, remains valid for

$$0 \leq \sigma = \theta_d/\delta^3 \leq 1/4.$$

At $s = 1/4$, equality in the puncture profile holds precisely when $A = C \setminus K$ and K is the union of two cosets of a finite subgroup L , whose difference has order greater than two in G/L . At $\sigma = 1/4$, equality in the oriented profile holds precisely for these configurations with the additional condition

$$K \cap (K + d) = \emptyset.$$

The strict outer hypothesis $\varepsilon < 1/4$ cannot be replaced by $\varepsilon \leq 1/4$ in either canonical repair assertion.

Proof. If $s = 1/4$, the construction in Theorem II.299.4 still gives $x \leq (2/3)\delta^3 < \delta/6$, so K is nonempty. The function $z(y)$ is continuous at zero, has $z(0) = 1/4$, and is strictly decreasing. For $y > 0$ its values are below $1/4$, so the derivative calculation (II.299.15) applies and is strictly negative. Continuity of F at zero therefore gives $F(x) \leq F(0)$, with strict inequality for $x > 0$. The inner coset theorem is valid at $\zeta = 1/4$, and its supplied hole has

$$|D| \leq \frac{3}{2}|K| < \frac{|C|}{2},$$

so properness is unchanged. The rest of the minimization argument applies without modification.

For the oriented assertion, $\theta_d \leq \delta^3/4 < \delta/16$ still forces $d \in H$. The endpoint puncture bound gives $e = \Delta_C(A) \leq \delta/2$. If a minimizing hole were preserved by d , the same intersection argument would give $z_d \geq \delta - e \geq \delta/2$, and consequently

$$\theta_d \geq z_d^2 \geq \frac{\delta^2}{4} > \frac{\delta^3}{4} \geq \theta_d,$$

a contradiction. Thus every minimizing hole is oriented correctly.

Suppose equality holds with $s = 1/4$. Strict decrease of F away from zero forces $x = 0$, and the complement identity gives $\zeta = s = 1/4$. The canonical inner coset must attain its envelope. The endpoint classification in Theorem II.298.5 says exactly that K is the stated union of two L -cosets.

To prove the converse, write $|L| = \ell$, so $|K| = 2\ell$. Every subgroup coset E has distance at least ℓ from K . Indeed, if E meets only one of the two L -cosets, their intersection has size $i \leq \ell$ and

$$|K \triangle E| = 2\ell + |E| - 2i \geq 2\ell - i \geq \ell.$$

If it meets both, each intersection has the same size $i = |L \cap L_E|$. When $i < \ell$, the subgroup index gives $i \leq \ell/2$, while $|E| \geq 2i$, so $|K \triangle E| = 2\ell + |E| - 4i \geq \ell$. When $i = \ell$, E contains both full L -cosets. Its index over L is at least three: index two would make their quotient difference have order two. Hence the same distance is at least ℓ . A coset meeting neither is farther still. The two constituent L -cosets attain distance ℓ , proving the minimum. Thus $\Delta_C(A) = \ell/m = \delta/2$ as required.

For oriented equality at $\sigma = 1/4$, the monotone chain of the proof of Theorem II.299.8 forces $\eta = \theta_d$ and endpoint puncture equality. With $x = 0$, (II.299.19) makes this additional equality equivalent to $r_K(d) = 0$. This is the displayed disjointness condition. The converse follows from the same identity.

Finally, let A be the union of two cosets of a nontrivial finite subgroup H , with quotient difference of order greater than two. Then $\varepsilon = 1/4$, and a canonical outer coset C is one of the two constituents. Here

$$\delta = 1/2, \quad \eta = 0, \quad K = C \setminus A = \emptyset.$$

The proposed puncture bound would require zero edits, although $A \neq C \setminus D$ for every hole $D \subset C$. Taking a translation between the two constituents gives $\theta_d = 0$ as well, so the oriented conclusion also fails at the outer endpoint. For a concrete finite cyclic instance, take $H = \{0, 3\} \subset \mathbb{Z}/6\mathbb{Z}$ and $A = \{0, 1, 3, 4\}$. $\square$

Example II.299.12 (Attainment at the inner endpoint). Let $G = \mathbb{Z}/8\mathbb{Z}$, $C = G$, $K = \{0, 1\}$, and $A = C \setminus K$. Then

$$m = 6, \quad \delta = \frac{1}{3}, \quad \varepsilon = \frac{25}{108}, \quad \eta = \frac{1}{108}, \quad s = \frac{1}{4}, \quad \Delta_C(A) = \frac{1}{6} = \delta\tau(1/4).$$

The singleton holes $\{0\}$ and $\{1\}$ both attain the minimum. For $d = 2$, $K \cap (K + d)$ is empty, so $\theta_d = \eta$, $\sigma = 1/4$, and the oriented endpoint is attained as well.

II.300 Exact ordered defects and endpoint reflection

Source 66, delivered Section 2.

Definition II.300.1. For a finite real set of size m , put

$$D(A) = \frac{M_m - E(A)}{4}. \quad (\text{II.300.1})$$

For the empty set, use $M_0 = E(\emptyset) = D(\emptyset) = 0$.

The following identity explains why the factor four is natural. It also proves that $D(A)$ is a nonnegative integer, which will matter when separating successive energy levels.

Lemma II.300.2 (Ordered-defect identity). *For every finite $A \subset \mathbb{R}$,*

$$D(A) = \#\{x < y < z : x, y, z \in A, x + z - y \notin A\}. \quad (\text{II.300.2})$$

If $A = B \cup \{a\}$, $a > \max B$, then

$$D(A) = D(B) + e(B, a). \quad (\text{II.300.3})$$

Proof. Let $m = |A|$ and $n = m - 1 = |B|$. Among additive quadruples involving a , the quadruple (a, a, a, a) contributes one. Those with exactly two occurrences of a contribute $4n$: the two occurrences lie on opposite sides of the equation, and the remaining entries are equal. There are no quadruples with exactly three occurrences, and no quadruples with two occurrences on one side and none on the other, by maximality of a .

A quadruple with exactly one occurrence of a is obtained by choosing an ordered pair $z < x$ in B for which $y = a - x + z$ belongs to B , then placing a, z on one side and x, y on the other. There are four choices for the side and position of a . This count is valid also when $x = y$: there is then one ordered choice for x, y , not two. Consequently, with

$$T(B, a) = \#\{z < x : z, x \in B, a - x + z \in B\},$$

we have

$$E(A) - E(B) = 4m - 3 + 4T(B, a). \quad (\text{II.300.4})$$

Since $T(B, a) + e(B, a) = \binom{n}{2}$ and $M_m - M_{m-1} = 2m^2 - 2m + 1$, equation (II.300.3) follows.

Every triple counted by (II.300.2) either lies in B , or has largest point a and is exactly a failed pair in (F.71.5). Induction from sets of size at most two proves the identity. $\square$

Lemma II.300.3 (Reflection). *Let B be nonempty, let $a > \max B$, and let $C = a - B$. Then*

$$e(B, a) = \sigma(C). \quad (\text{II.300.5})$$

Proof. A pair $z < x$ in B gives the pair $a - z > a - x$ in C . The difference of these two positive numbers is $x - z$. Membership of $a - x + z$ in B is equivalent to membership of $x - z$ in C . This is a bijection between the two sets of failures. $\square$

If $c_1 < \dots < c_n$ are the positive elements of C , write

$$e_i = \#\{j < i : c_i - c_j \notin C\}, \quad \sigma(C) = \sum_{i=1}^n e_i. \quad (\text{II.300.6})$$

A positive difference smaller than c_i can belong only to $C_{i-1} = \{c_1, \dots, c_{i-1}\}$. Thus e_i is determined entirely by the prefix C_i . Equivalently, $\binom{n}{2} - \sigma(C)$ is the number of ordered positive Schur representations $u + v = w$ within C ; repeated summands $u = v$ are allowed and are counted once for that ordered choice.

II.301 Two exact reflections force an arithmetic prefix

Source 66, delivered Section 3.

Lemma II.301.1 (Consecutive zero defects). *Suppose $1 \leq i < n$ and $e_i = e_{i+1} = 0$. With $d = c_1$, one has*

$$C_{i+1} = \{d, 2d, \dots, (i+1)d\}.$$

Proof. For $i = 1$, the assertion is $c_2 - c_1 = c_1$, which follows immediately from $e_2 = 0$. Assume $i \geq 2$. Since $e_i = 0$, reflection $s \mapsto c_i - s$ permutes C_{i-1} . Since $e_{i+1} = 0$, reflection $s \mapsto c_{i+1} - s$ permutes C_i . Pairing the smallest and largest elements of C_i under the latter reflection gives

$$c_{i+1} = c_i + d. \quad (\text{II.301.1})$$

For $s \in C_{i-1}$, composing the two reflections shows that

$$s + d = c_{i+1} - (c_i - s) \in C_i.$$

Thus every point of C_i other than its largest point has a successor obtained by adding d . Starting at any point and repeatedly adding d must eventually reach c_i , since the set is finite. In particular, starting at d shows that c_i is an integer multiple of d and that all intermediate multiples occur. Starting at any other point shows that it belongs to the same chain. There are i points, so $c_i = id$; (II.301.1) finishes the proof. $\square$

Corollary II.301.2 (Zero defect). *A positive n -element set has $\sigma(C) = 0$ if and only if $C = \{d, 2d, \dots, nd\}$ for some $d > 0$. Consequently, $e(B, a) = 0$ if and only if $B \cup \{a\}$ is an arithmetic progression whose last point is a .*

Proof. The cases $n \leq 2$ are immediate. For $n \geq 3$, apply Lemma II.301.1 to the final pair of indices. The converse is immediate from closure of the relevant positive differences. Reflection then proves the endpoint assertion. $\square$

Lemma II.301.3 (A long initial arithmetic block). *Suppose $\sigma(C) \leq r$, and let L be the largest integer such that $c_j = jd$ for $1 \leq j \leq L$, where $d = c_1$. Then*

$$L \geq n - 2r. \quad (\text{II.301.2})$$

Proof. There is nothing to prove when $L = n$. Otherwise $e_{L+1} > 0$: zero would make reflection of $\{d, \dots, Ld\}$ about $c_{L+1}/2$ a permutation, forcing $c_{L+1} = (L+1)d$. After this first positive defect, no two consecutive e_i can vanish, by Lemma II.301.1 and maximality of L . A zero-positive pattern starting with a positive entry and containing no consecutive zeros has at most twice as many entries as positive entries. There are at most r positive entries, since each costs at least one in $\sum e_i$. Hence $n - L \leq 2r$. $\square$

II.302 The partition normal-form theorem

Source 66, delivered Section 4.

Theorem II.302.1 (Small Schur defect). *Let $r \geq 1$ be an integer and $n \geq 3r + 1$. For a positive ordered set $C = \{c_1 < \dots < c_n\}$, the following are equivalent:*

- (i) $\sigma(C) \leq r$.
- (ii) With $d = c_1$, there are nonnegative integers $0 \leq \lambda_1 \leq \dots \leq \lambda_n$ such that

$$c_i = d(i + \lambda_i), \quad \sum_i \lambda_i \leq r.$$

In this situation $\sigma(C) = \sum_i \lambda_i$. In particular, exact defect r gives precisely $p(r)$ classes up to positive dilation.

Proof. Assume (i). Let L be as in Lemma II.301.3; then $L \geq n - 2r \geq r + 1$. If $L = n$, all conclusions hold with zero shifts. We divide the remaining argument into three steps.

First, all points are integer multiples of d . Otherwise let c_i be the first point that is not. Every previous point is a multiple of d , and every difference $c_i - c_j$, $j < i$, is not. None belongs to the previous prefix, so $e_i = i - 1 \geq L > r$, a contradiction. Scale by d and henceforth take $d = 1$.

Second, all points are at most $2L + 1$. Suppose that i is the first index with $c_i \geq 2L + 2$. For every earlier index j , write $q_j = c_j - j$, the number of missing positive integers below c_j . All missing integers are at least $L + 1$. Since $c_j \leq 2L + 1$, two missing integers cannot sum to c_j . Among the $c_j - 1$ ordered decompositions $u + v = c_j$ into positive integers, exactly q_j fail through u and exactly q_j fail through v , and these failure classes are disjoint. Thus

$$(j - 1) - e_j = c_j - 1 - 2q_j, \quad e_j = q_j. \quad (\text{II.302.1})$$

Let $s = i - 1 - L$ be the number of earlier points beyond the initial block. Every such point has $q_j \geq 1$, hence $\sum_{j < i} e_j \geq s$. For each $t \in \{1, \dots, L\}$, the difference $c_i - t$ is larger than L . At most s of these L distinct differences can belong to the earlier prefix. Therefore $e_i \geq L - s$, and

$$\sigma(C) \geq \sum_{j < i} e_j + e_i \geq s + (L - s) = L > r,$$

again a contradiction.

Third, the shifts have exactly the claimed total. Now every target is at most $2L + 1$, so the same disjoint-hole count (II.302.1) applies to every index. Set $\lambda_i = c_i - i$. These are nonnegative integers, and they are nondecreasing because $c_{i+1} \geq c_i + 1$. Their sum is $\sum_i e_i = \sigma(C)$.

Conversely, assume (ii). At most r shifts are positive, so the first $n - r$ points are $1, \dots, n - r$ after scaling. Also $c_n \leq n + r$. Two missing positive integers have sum at least $2(n - r + 1)$, which is larger than $n + r$ when $n \geq 3r + 1$. Thus the disjoint-hole count applies again and gives $e_i = \lambda_i$ for every i .

Removing the zero shifts gives a nondecreasing list of positive integers summing to the exact defect. Conversely, every partition of r has at most $r < n$ entries, and padding it with leading zeros constructs a valid set. The scale is fixed by $c_1 = d$, so different lists give different positive-dilation classes. $\square$

[write] **7 October 2026; sources 60–69.** *Against source 57* (Section II.291; checked at the write). Source 57's Theorem II.292.3 has the same conclusion for $n \geq 5s - 1$ (positions among the last s); $3r + 1 \leq 5r - 1$ with equality at $r = 1$, so this theorem's threshold is better for $r \geq 2$, except that source 57 proves the optimal value 6 at $r = 2$ (Theorem II.292.4). With Proposition II.302.4, the optimal threshold satisfies $2r + 2 \leq N(r) \leq 3r + 1$; $N(1) = 4$ and $N(2) = 6$ attain the lower end, as this source's Conjecture V.249.1 predicts.

Corollary II.302.2 (A finite set of local obstructions). *Under the hypotheses of Theorem II.302.1, C contains $\{d, 2d, \dots, (n-r)d\}$ and is contained in $\{d, 2d, \dots, (n+r)d\}$. The number of classes with defect at most r is $\sum_{j=0}^r p(j)$, where $p(0) = 1$.*

Example II.302.3 (Defects one, two, and three). Suppressing leading zeros in the shift vector, the possibilities are

r	positive shifts
1	(1)
2	(2), (1, 1)
3	(3), (1, 2), (1, 1, 1).

For example, defect two gives either $\{1, \dots, n-1, n+2\}$ or $\{1, \dots, n-2, n, n+1\}$, up to dilation. These two sets produce different boundary-hole descriptions after reflection; the partition parameter need not be unchanged by that reflection.

Proposition II.302.4 (A genuine size obstruction). *For every $r \geq 1$, the set*

$$C_r = [1, 2r+2]_{\mathbb{Z}} \setminus \{r+1\}$$

has $2r+1$ points and $\sigma(C_r) = r$, but its shift sum is $r+1$. Thus an unrestricted theorem of the form of Theorem II.302.1 cannot start at $n = 2r+1$.

Proof. Every target above the missing point has one hole below it, so the uncorrected shift sum is $r+1$. The only target that is a sum of two holes is $2r+2 = (r+1) + (r+1)$. At this target, the two failure classes in the inclusion–exclusion count intersect in one ordered choice. That reduces the Schur defect by one. All other targets have no such intersection. Therefore $\sigma(C_r) = r$. $\square$

The obstruction identifies the phenomenon that the size hypothesis prevents: *two holes can repair a Schur deficit by adding to an existing point*. The bound $3r+1$ is proved, not asserted optimal. The interval between $2r+2$ and $3r+1$ is a concrete further problem.

II.303 Boundary-hole normal forms for endpoint extensions

Source 66, delivered Section 5.

Definition II.303.1 (Hole weight). For a finite positive integer set $H = \{h_1 < \dots < h_s\}$, define

$$\text{wt}(H) = \sum_{j=1}^s (h_j - j + 1) = \sum_{h \in H} h - \binom{s}{2}. \quad (\text{II.303.1})$$

The empty hole set has weight zero. For nonempty H and $N > \max H$ (or $H = \emptyset$ and $N \geq 0$), put

$$A_H(N) = [0, N]_{\mathbb{Z}} \setminus H. \quad (\text{II.303.2})$$

The numbers $\mu_j = h_j - j + 1$ are positive and nondecreasing. Conversely, any positive nondecreasing list $\mu_1, \dots, \mu_s$ gives $h_j = \mu_j + j - 1$. This is a bijection between hole sets of weight r and partitions of r . It also gives

$$s \leq r, \quad h_s \leq r. \quad (\text{II.303.3})$$

The second assertion follows directly from $\sum_{j < s} h_j \geq \binom{s}{2}$.

Theorem II.303.2 (Endpoint normal forms). *Let $A = B \cup \{a\} \subset \mathbb{R}$, where $a > \max B$, let $|B| = n$, and suppose $e(B, a) = r \geq 1$ and $n \geq 3r + 1$. Then there is a unique normalization by translation and positive dilation for which*

$$A = A_H(N), \quad \text{wt}(H) = r, \quad N = n + |H|. \quad (\text{II.303.4})$$

In this normalization the gap between the last two points is one, and $H \subset [1, r]_{\mathbb{Z}}$. Conversely, every hole set of weight r , with $N = n + |H|$ and $n \geq 3r + 1$, produces endpoint defect r . Consequently, for fixed n and r in this range, there are exactly $p(r)$ oriented affine classes.

Proof. Translate so that $\min B = 0$, and let $d = a - \max B$. Reflection gives $C = a - B$ with $c_1 = d$. By Theorem II.302.1, after dividing by d we may write $c_i = i + \lambda_i$. Let

$$N = c_n = n + \lambda_n,$$

so $a = N$. The reflected endpoint set is $\{N\} \cup (N - C)$, an integer interval with a hole set H . Its number of holes is $s = N - n = \lambda_n$.

Let t be the number of positive shifts. The initial block of C has length $n - t$. Every hole of C is at least $n - t + 1$, so after reflection every hole of A is at most

$$N - (n - t + 1) = s + t - 1 \leq \sum_i \lambda_i = r.$$

The inequality uses the $t - 1$ other positive shifts, each at least one. Finally, summing the points in the reflected set gives

$$\begin{aligned} \sum_i \lambda_i &= \sum_{c \in C} c - \frac{n(n+1)}{2} \\ &= nN - \left(\frac{N(N-1)}{2} - \sum_{h \in H} h \right) - \frac{n(n+1)}{2} \\ &= \sum_{h \in H} h - \binom{s}{2} = \text{wt}(H). \end{aligned}$$

This proves the asserted form. Translation is fixed by the minimum, and dilation by the final gap d , so the normalization is unique.

For the converse, take H of weight r and $N = n + s$. Since $N \geq 3r + 1 + s > 2 \max H$, the last gap is one. To count failed pairs at endpoint N , fix a missing completion $h \in H$ and a present point z with $0 \leq z < h$. The other point is necessarily

$$x = N - h + z.$$

It lies in $[0, N - 1]_{\mathbb{Z}} \setminus H$: indeed, $x \geq N - \max H > \max H$. Conversely, every failed pair is obtained this way. There are $h_j - (j - 1)$ choices for z at the hole h_j . Their sum is $\text{wt}(H) = r$. The partition bijection and uniqueness of normalization finish the proof. $\square$

Corollary II.303.3 (Full-set containment from the local defect). *Under the hypotheses of Theorem II.303.2, A is contained in an arithmetic progression of at most $m + r$ points and contains an arithmetic progression of at least $m - r$ points, where $m = n + 1$.*

Proof. The containing interval has $N + 1 = m + s \leq m + r$ points. The terminal interval $[\max H + 1, N]_{\mathbb{Z}}$ has $N - \max H = m - 1 + s - \max H \geq m - r$ points, since $s \geq 1$ and $\max H \leq r$. $\square$

Remark II.303.4 (An orientation-sensitive parameter). The reflected set of an endpoint model has the same total energy but usually a different last-endpoint defect. For example, $[0, m]_{\mathbb{Z}} \setminus \{1\}$ has final defect one in the theorem's range, whereas its reflection $[0, m]_{\mathbb{Z}} \setminus \{m - 1\}$ has final defect $m - 2$. Thus an endpoint classification must not be silently substituted for a classification under a total-energy hypothesis.

II.304 An exact energy law for arbitrary boundary holes

Source 66, delivered Section 6.

The normal-form theorem reduces an infinite real-variable problem to finitely many hole sets. The next formula computes the entire energy of each such template, not only the local defect.

For $H = \{h_1 < \cdots < h_s\}$ define

$$\mathcal{P}(H) = \{(h, x) : h \in H, 0 \leq x < h, x \notin H\}. \quad (\text{II.304.1})$$

It has $\text{wt}(H)$ elements. For $(h, x) \in \mathcal{P}(H)$ let

$$R_H(h, x) = \#\{(u, v) \in H^2 : v > h, v - u = h - x\}. \quad (\text{II.304.2})$$

Write $j_H(h) = \#(H \cap [1, h]_{\mathbb{Z}})$ and set

$$\kappa(H) = \sum_{(h, x) \in \mathcal{P}(H)} (h + s - j_H(h) - \#(H \cap [1, x]_{\mathbb{Z}}) - R_H(h, x)). \quad (\text{II.304.3})$$

An interval with upper endpoint zero in the counting term is empty. Use $\kappa(\emptyset) = 0$.

Theorem II.304.1 (Exact boundary-hole energy). *Let H be a finite positive integer set, let $r = \text{wt}(H)$, and suppose $N \geq 2 \max H$ when $H \neq \emptyset$. If $m = |A_H(N)| = N + 1 - |H|$, then*

$$D(A_H(N)) = r(m - 1) - \kappa(H), \quad (\text{II.304.4})$$

$$E(A_H(N)) = M_m - 4r(m - 1) + 4\kappa(H). \quad (\text{II.304.5})$$

For $H = \emptyset$, the same formulas hold with $r = \kappa = 0$.

Proof. By Lemma II.300.2, a failed triple $x < y < z$ has a missing completion $h = x + z - y$. This completion lies strictly between x and z , hence belongs to H . Conversely, choose $(h, x) \in \mathcal{P}(H)$ and set $t = h - x > 0$. A corresponding triple is determined by

$$z \in [h + 1, N]_{\mathbb{Z}}, \quad y = z - t,$$

provided both z and y are present. The inequality $z > h$ gives $y > x$, so there is no additional ordering constraint.

There are initially $N - h$ choices for z . Exactly $s - j_H(h)$ are excluded because $z \in H$. Choices excluded because $y \in H$ are in bijection with $u \in H$, $u > x$, by $z = u + h - x$. All of these satisfy $z \leq N$: the hypothesis gives $u + h - x \leq 2 \max H \leq N$. Their number is $s - \#(H \cap [1, x]_{\mathbb{Z}})$. The

overlap of the two exclusion classes is precisely $R_H(h, x)$. Therefore the number of valid triples for this pair is

$$N - h - (s - j_H(h)) - (s - \#(H \cap [1, x]_{\mathbb{Z}})) + R_H(h, x).$$

Summing over the r pairs, and using $N = m - 1 + s$, gives (II.304.4). The energy formula follows from the definition of D . $\square$

Remark II.304.2 (Why the overlap term is indispensable). When two holes are separated by $h - x$, the same choice of z is excluded twice by the two single-hole tests. Omitting R_H therefore changes the constant term of the energy. For example, two distinct weight-four partitions already have the same value of κ ; the energy law does not imply that energy alone determines the partition. The local normal form is strictly more informative.

II.304.1 The first values

Table II.304.1 lists the complete spectrum through weight four. The formula for D uses $n = m - 1$ to make the endpoint interpretation visible. Every entry is an exact identity in the stated long-interval range.

r	H	$\kappa(H)$	$D(A_H(N))$
1	$\{1\}$	1	$n - 1$
2	$\{2\}$	4	$2n - 4$
	$\{1, 2\}$	3	$2n - 3$
3	$\{3\}$	9	$3n - 9$
	$\{1, 3\}$	7	$3n - 7$
	$\{1, 2, 3\}$	6	$3n - 6$
4	$\{4\}$	16	$4n - 16$
	$\{1, 4\}$	12	$4n - 12$
	$\{2, 3\}$	11	$4n - 11$
	$\{1, 2, 4\}$	10	$4n - 10$
	$\{1, 2, 3, 4\}$	10	$4n - 10$

Table II.304.1: Partition-indexed endpoint templates. The repeated constant at weight four is an actual energy coincidence, not a rounding effect.

For instance, $H = \{2\}$ has pairs $(2, 0), (2, 1)$ and contributes two summands equal to two in (II.304.3), so $\kappa = 4$. For $H = \{1, 2\}$, the only pairs are $(1, 0), (2, 0)$. Their contributions are respectively one and two, giving $\kappa = 3$. At $H = \{1, 3\}$, the contributions of $(1, 0), (3, 0), (3, 2)$ are 2, 3, 2, hence $\kappa = 7$. These are the constants used in the energy classification, so none depends on numerical approximation.

II.305 A sharp all-defect endpoint envelope

Source 66, delivered Section 7.

Lemma II.305.1 (The largest boundary constant). *If $\text{wt}(H) = r \geq 1$, then*

$$\kappa(H) \leq r^2, \tag{II.305.1}$$

with equality if and only if $H = \{r\}$.

Proof. There are r summands in (II.304.3). For a pair with $h = h_j$, its summand is at most

$$h_j + s - j \leq h_s \leq r.$$

The middle inequality follows because the $s - j$ later holes are strictly increasing integers. This proves (II.305.1).

If equality holds, every summand equals r and $h_s = r$. Equality in the elementary inequality $h_s \leq \text{wt}(H)$ forces $h_j = j$ for $j < s$. If $s \geq 2$, consider the pair $(1, 0)$. When $r > s$, its contribution is at most $s < r$. When $r = s$, the hole set is $\{1, \dots, r\}$, and $R_H(1, 0) \geq 1$, so that contribution is again smaller than r . Both contradict equality. Therefore $s = 1$ and $H = \{r\}$. Conversely, the r pairs (r, x) , $0 \leq x < r$, each contribute r for that hole set. $\square$

Theorem II.305.2 (Sharp endpoint-sensitive energy bound). *Let $A = B \cup \{a\} \subset \mathbb{R}$, $a > \max B$, and $|A| = m$. Suppose $e(B, a) = r \geq 1$ and $m - 1 \geq 3r + 1$. Then*

$$D(A) \geq r(m - 1 - r), \quad (\text{II.305.2})$$

$$E(A) \leq M_m - 4r(m - 1 - r). \quad (\text{II.305.3})$$

Equality in either bound holds if and only if

$$A = b + d([0, m]_{\mathbb{Z}} \setminus \{r\}), \quad d > 0. \quad (\text{II.305.4})$$

The coefficient and the constant term are therefore sharp for every allowed r and m .

Proof. Theorem II.303.2 gives $A_H(N)$ with $\text{wt}(H) = r$, $N = m - 1 + |H|$, and $\max H \leq r$. Its size hypothesis implies $N > 2 \max H$, so Theorem II.304.1 applies. Use Lemma II.305.1; its equality case gives exactly (II.305.4). Conversely, that model has endpoint defect r by Theorem II.303.2, and $\kappa(\{r\}) = r^2$. $\square$

Corollary II.305.3 (The smallest three endpoint defects). *Put $n = m - 1$. For final endpoint defect one and $n \geq 4$, the only normal form is $A_{\{1\}}(n + 1)$ and $D = n - 1$. For defect two and $n \geq 7$, the possibilities are*

$$A_{\{2\}}(n + 1), \quad D = 2n - 4; \quad A_{\{1,2\}}(n + 2), \quad D = 2n - 3.$$

For defect three and $n \geq 10$, the possibilities and their defects are

H	D
$\{3\}$	$3n - 9$
$\{1, 3\}$	$3n - 7$
$\{1, 2, 3\}$	$3n - 6$.

In each line $N = n + |H|$.

The significance of Theorem II.305.2 is that the entire endpoint spectrum has been replaced by a finite partition problem. It is not merely a bound obtained by discarding a few exceptional points. Nevertheless, it controls a local endpoint parameter; a separate argument is needed to infer a normal form from a large total energy. We give that argument for four levels next.

II.306 The known first and second levels, with complete proofs

Source 66, delivered Section 8.

The statements in this section, including the third-level lower bound, are already recorded in the source-45 repository review [266]. Their proofs are included to isolate the genuinely new equality analysis in Sections II.307 and II.308.

For $m \geq 5$, define the oriented models

$$P_m = [0, m-1]_{\mathbb{Z}}, \quad L_m = [0, m]_{\mathbb{Z}} \setminus \{1\}, \quad R_m = [0, m]_{\mathbb{Z}} \setminus \{m-1\}. \quad (\text{II.306.1})$$

They have defects $0, m-2, m-2$, respectively.

Lemma II.306.1 (Extending an arithmetic progression). *Let $B = [0, n-1]_{\mathbb{Z}}$ and $a > n-1$. If $a \notin \mathbb{Z}$, then $e(B, a) = \binom{n}{2}$. If $a = n+t$ with $t \in \mathbb{Z}_{\geq 0}$, then*

$$e(B, a) = \sum_{j=1}^{\min(t, n-1)} (n-j). \quad (\text{II.306.2})$$

In particular, for $0 \leq t \leq n-1$, it equals $tn - t(t+1)/2$.

Proof. For a pair $z < x$, write $j = x - z$. There are $n-j$ pairs of difference j . A noninteger endpoint makes every completion nonintegral. At endpoint $n+t$, the completion belongs to B exactly when $n+t-j \leq n-1$, or $j \geq t+1$. The remaining pairs are precisely those counted in (II.306.2). $\square$

Proposition II.306.2 (First gap and second equality cases). *For a real m -element set with $m \geq 3$, either A is an arithmetic progression or*

$$D(A) \geq m-2. \quad (\text{II.306.3})$$

For $m \geq 5$, equality in (II.306.3) holds if and only if A is an oriented affine copy of L_m or R_m .

Proof. For three points, the sole ordered triple has its completion in the set exactly when the three points form an arithmetic progression. Thus the defect is zero or one. Proceed by induction, writing $A = B \cup \{a\}$ and $n = |B| = m-1$.

If B is an arithmetic progression, normalize it to $[0, n-1]_{\mathbb{Z}}$. Lemma II.306.1 gives either the next arithmetic point, or a defect at least $n-1 = m-2$. If B is not an arithmetic progression, induction gives $D(B) \geq n-2$ and Corollary II.301.2 gives $e(B, a) \geq 1$. The recurrence proves the gap.

Now let $m \geq 5$ and $D(A) = m-2 = n-1$. If B is an arithmetic progression, (II.306.2) shows that the only endpoint is $a = n+1$; the noninteger value $\binom{n}{2}$ and all integer choices $t \geq 2$ are larger. This is R_m . If B is not an arithmetic progression, the preceding lower bounds force $e(B, a) = 1$. Here $n \geq 4$, so Corollary II.305.3 gives L_m . The asserted defects of L_m, R_m follow from Theorem II.304.1 and reflection. $\square$

Lemma II.306.3 (Extending either second-level model). *Let $n \geq 4$, let $B = L_n$ or R_n , and let $a > n$. Unless $B = L_n$ and $a = n+1$, one has*

$$e(B, a) \geq n-2. \quad (\text{II.306.4})$$

For later use, the following more specific statements hold:

B	a	Endpoint defect or lower bound
L_n	$n + 1$	1
L_n	$n + 2$	$n - 2$
L_n	integer $a \geq n + 3$	at least $2n - 4$
R_n	$n + 1$	$n - 2$
R_n	$n + 2$	at least $2n - 5$
R_n	integer $a \geq n + 3$	at least $2n - 4$
Either	noninteger a	$\binom{n}{2}$

Proof. For any positive difference j of B , every pair with that difference fails exactly when $a - j \notin B$; its contribution is the multiplicity $r_B(j) = \#\{z \in B : z + j \in B\}$. Both L_n and R_n satisfy

$$r_B(1) = r_B(2) = n - 2.$$

For L_n at $a = n + 1$, the only failing difference is n , which has multiplicity one. At $a = n + 2$, the only failing difference is one, with multiplicity $n - 2$; the missing completion one would require difference $n + 1$, which is unavailable. At $a \geq n + 3$, differences one and two both fail.

For R_n at $a = n + 1$, the only failing difference is two, again with multiplicity $n - 2$. At $a = n + 2$, differences one and three fail; their multiplicities are $n - 2$ and $n - 3$. At $a \geq n + 3$, differences one and two fail. Noninteger endpoints make every completion noninteger. These observations prove the table and the lower bound. $\square$

Proposition II.306.4 (Third-level exclusion bound). *Let $m \geq 5$. If $A \subset \mathbb{R}$ has m points and is not an oriented affine copy of P_m, L_m, R_m , then*

$$D(A) \geq 2m - 6. \quad (\text{II.306.5})$$

Proof. At $m = 5$, Proposition II.306.2 and its equality classification exclude defects zero through three except the listed sets. Integrality therefore gives $D(A) \geq 4 = 2m - 6$.

For the induction step, let $m \geq 6$, put $n = m - 1$, and remove the last point. If B is an arithmetic progression, the cases $t = 0, 1$ in Lemma II.306.1 give excluded sets. Every other endpoint has defect at least $2n - 3$, including noninteger endpoints since $n \geq 5$; this is larger than $2n - 4 = 2m - 6$.

If B is a second-level model, its defect is $n - 2$. Lemma II.306.3 adds at least $n - 2$, except for the extension of L_n to L_m , which is excluded.

In the remaining case, induction gives $D(B) \geq 2n - 6$. Endpoint defect zero is impossible by Corollary II.301.2; endpoint defect one would make A a copy of L_m by Corollary II.305.3, also excluded. Hence $e(B, a) \geq 2$, and the recurrence gives $D(A) \geq 2n - 4$. $\square$

II.307 Complete classification of the third energy level

Source 66, delivered Section 9.

Define three oriented templates

$$\begin{aligned} J_m^- &= [0, m]_{\mathbb{Z}} \setminus \{2\}, & J_m^+ &= [0, m]_{\mathbb{Z}} \setminus \{m - 2\}, \\ C_m &= [0, m + 1]_{\mathbb{Z}} \setminus \{1, m\}. \end{aligned} \quad (\text{II.307.1})$$

The last model has one hole near each endpoint. It cannot be found by looking only for the reflected single-hole construction.

Theorem II.307.1 (All third-level extremizers). *Let $m \geq 8$. A real m -element set satisfies*

$$D(A) = 2m - 6 \quad \text{equivalently} \quad E(A) = M_m - 8(m - 3) \quad (\text{II.307.2})$$

if and only if it is an oriented affine copy of one of J_m^- , J_m^+ , C_m .

Proof. The defect of J_m^- is $2(m - 1) - 4 = 2m - 6$ by Theorem II.304.1; reflection gives the same for J_m^+ . The model C_m is obtained by extending L_{m-1} at endpoint $m + 1$. The defect of the prefix is $m - 3$, and Lemma II.306.3 gives the same endpoint defect $m - 3$. Thus all three models attain (II.307.2).

Conversely, put $n = m - 1 \geq 7$ and write $A = B \cup \{a\}$ with the last point removed. We distinguish the three prefix types used in Proposition II.306.4.

Arithmetic prefix. The extension is either an excluded first- or second-level model, or has defect at least $2n - 3$. The target defect is $2n - 4$, so this case gives no equality.

Second-level prefix. Normalize B to L_n or R_n . Since $D(B) = n - 2$, equality requires endpoint defect exactly $n - 2$. For L_n , the endpoint $n + 1$ has defect one, and the endpoint $n + 2$ has defect $n - 2$, yielding C_m . All later integer endpoints have defect at least $2n - 4 > n - 2$. Noninteger endpoints have defect $\binom{n}{2} > n - 2$. For R_n , endpoint $n + 1$ has defect $n - 2$ and yields J_m^+ . The endpoint $n + 2$ has defect at least $2n - 5 > n - 2$ because $n \geq 7$; all later integer endpoints and noninteger endpoints are also too large. This exhausts the second-level prefixes.

Any other prefix. Proposition II.306.4 gives $D(B) \geq 2n - 6$, so the target forces $e(B, a) \leq 2$. Endpoint defects zero and one would make A a first- or second-level model, and are impossible. Thus $e(B, a) = 2$. Since $n \geq 7$, Corollary II.305.3 applies: the two normal forms have total defects $2n - 4$ and $2n - 3$. Only $A_{\{2\}}(n + 1) = J_m^-$ has the required defect. $\square$

[write] **7 October 2026; sources 60–69.** *Against source 57* (checked at the write). Source 57's Theorem II.292.7 classifies the same level for $m \geq 7$, a better cutoff, as the affine images of $J_m = \{0, \dots, m\} \setminus \{2\}$ and $K_m = \{0, 2, 3, \dots, m - 1, m + 1\}$: this theorem's J_m^- and J_m^+ are J_m and its reflection, and C_m is K_m , which is symmetric. Both answer source 45's Question V.166.1.

Remark II.307.2 (The cutoff is not cosmetic). At $m = 6$, the set $[0, 6]_{\mathbb{Z}} \setminus \{3\}$ also has defect six, which equals $2m - 6$, but is not one of the three displayed templates. The theorem makes no claim of optimality for the sufficient cutoff $m \geq 8$. The example explains why a small-cardinality qualification cannot simply be omitted. The inclusion of C_m is essential even when m is arbitrarily large.

II.308 The fourth energy level and the eight-template theorem

Source 66, delivered Section 10.

Put

$$T_m^- = [0, m + 1]_{\mathbb{Z}} \setminus \{1, 2\}, \quad T_m^+ = [0, m + 1]_{\mathbb{Z}} \setminus \{m - 1, m\}. \quad (\text{II.308.1})$$

These are reflected copies of one another.

Theorem II.308.1 (All fourth-level extremizers). *Let $m \geq 11$. A real m -element set satisfies*

$$D(A) = 2m - 5 \quad \text{equivalently} \quad E(A) = M_m - 8m + 20 \quad (\text{II.308.2})$$

if and only if it is an oriented affine copy of T_m^- or T_m^+ .

Proof. The boundary-hole formula gives $D(T_m^-) = 2(m-1) - 3 = 2m - 5$, and reflection gives the same for T_m^+ . For the converse, put $n = m - 1 \geq 10$ and remove the last point.

Arithmetic prefix. With $B = [0, n-1]_{\mathbb{Z}}$, the target defect is $2n - 3$. Lemma II.306.1 gives that value exactly when $a = n + 2$. Values $a = n, n + 1$ give smaller defects; $a = n + t$ with $t \geq 3$ gives at least $3n - 6 > 2n - 3$; and $\binom{n}{2} > 2n - 3$ excludes noninteger endpoints. The resulting set is T_m^+ .

Second-level prefix. Here $D(B) = n - 2$, so the required endpoint defect is $n - 1$. The table in Lemma II.306.3 excludes it. For L_n , the first two possible integer endpoints have defects $1, n - 2$, and all later ones have at least $2n - 4 > n - 1$. For R_n , the first has defect $n - 2$, the second at least $2n - 5 > n - 1$, and later ones at least $2n - 4$. Noninteger endpoints are again larger. Thus no second-level prefix contributes.

Any other prefix. By Proposition II.306.4, $D(B) \geq 2n - 6$, and therefore $e(B, a) \leq 3$. Defects zero and one are excluded as before. Defect two gives the two models of Corollary II.305.3, and only the model with holes $\{1, 2\}$ has defect $2n - 3$. It is T_m^- . For endpoint defect three, that same corollary is valid since $n \geq 10$. The three possible total defects are

$$3n - 9, \quad 3n - 7, \quad 3n - 6.$$

All exceed $2n - 3$ when $n \geq 10$. This eliminates the final case. $\square$

Theorem II.308.2 (The first four levels). *For every $m \geq 11$, the four largest distinct additive energies of real m -element sets are the values in Table II.308.1, with exactly the indicated oriented affine models. In particular, every other m -element set has*

$$D(A) \geq 2m - 4, \quad E(A) \leq M_m - 8m + 16. \quad (\text{II.308.3})$$

The bound (II.308.3) is only an exclusion bound; it is not asserted to be the fifth energy level.

Proof. The maximum and its equality case follow from the ordered-defect identity and Corollary II.301.2, by removing the last point, or directly from Proposition II.306.2. The first gap and its equality classification give the second level. Proposition II.306.4 excludes every intervening defect before the third level, and Theorem II.307.1 classifies its extremizers. Theorem II.308.1 classifies the next integer value of the defect. Each displayed model exists for all $m \geq 11$, and the four defect values are distinct. All remaining defects are larger integers, giving (II.308.3). $\square$

$D(A)$	$E(A)$	Normalized oriented models
0	M_m	$[0, m-1]_{\mathbb{Z}}$
$m - 2$	$M_m - 4m + 8$	$[0, m]_{\mathbb{Z}} \setminus \{1\}, [0, m]_{\mathbb{Z}} \setminus \{m-1\}$
$2m - 6$	$M_m - 8m + 24$	$[0, m]_{\mathbb{Z}} \setminus \{2\}, [0, m]_{\mathbb{Z}} \setminus \{m-2\},$ $[0, m+1]_{\mathbb{Z}} \setminus \{1, m\}$
$2m - 5$	$M_m - 8m + 20$	$[0, m+1]_{\mathbb{Z}} \setminus \{1, 2\}, [0, m+1]_{\mathbb{Z}} \setminus \{m-1, m\}$

Table II.308.1: The first four energy levels for $m \geq 11$. The models are listed up to translation and *positive* dilation. Allowing reflection identifies the displayed pairs.

Example II.308.3 (The smallest cardinality covered by all four levels). At $m = 11$, $M_m = 891$, and the four energies are

$$891, \quad 855, \quad 827, \quad 823.$$

The third level includes the two one-hole sets $[0, 11]_{\mathbb{Z}} \setminus \{2\}$ and $[0, 11]_{\mathbb{Z}} \setminus \{9\}$, and the two-sided-hole set $[0, 12]_{\mathbb{Z}} \setminus \{1, 11\}$. The fourth level consists of $[0, 12]_{\mathbb{Z}} \setminus \{1, 2\}$ and its reflection. These numbers are evaluations of exact formulas, not estimates.

II.309 Inverse consequences and the boundary of the Ramsey application

Source 66, delivered Section 11.

II.309.1 Entire-set recovery, rather than a proportional subset

Corollary II.309.1 (A sharp high-energy recovery window). *Let $A \subset \mathbb{R}$ have $m \geq 11$ elements and suppose*

$$E(A) \geq \frac{2m^3 + m}{3} - 8m + 20. \quad (\text{II.309.1})$$

Then A is one of the eight oriented templates in Table II.308.1. It lies in an arithmetic progression of at most $m + 2$ points, contains an arithmetic progression of at least $m - 2$ points, and satisfies

$$|A - A| \leq 2m + 3. \quad (\text{II.309.2})$$

The upper bound on the containing progression and the difference-set bound are attained within this window.

Proof. The hypothesis is $D(A) \leq 2m - 5$, so Theorem II.308.2 gives the list. Every model is contained in $[0, m + 1]_{\mathbb{Z}}$. The consecutive blocks giving the inner progression have lengths respectively m , $m - 1$, at least $m - 2$, and $m - 1$ in the four rows of the table. All differences belong to $[-(m + 1), m + 1]_{\mathbb{Z}}$ after normalization, proving (II.309.2).

For T_m^- , the interval $[3, m + 1]_{\mathbb{Z}}$ supplies differences $1, \dots, m - 2$, and subtraction from zero supplies $3, \dots, m + 1$. These cover all positive differences through $m + 1$, so $|T_m^- - T_m^-| = 2m + 3$. Moreover, it has diameter $m + 1$ and contains consecutive integer points. Any real arithmetic progression containing it has step at most one: the difference one must be an integer multiple of the step. Such a progression has at least $m + 2$ points. $\square$

The comparison with Gowers's Theorem 7.2 is quantitative but local. Here the energy-density assumption is

$$\frac{E(A)}{m^3} \geq \frac{2}{3} - \frac{23}{3m^2} + \frac{20}{m^3}.$$

It lies within $O(m)$ of the ordered maximum and approaches $2/3$, whereas Gowers's extraction works at every fixed positive c_0 . The present corollary retains *all* points, gives dimension one, and loses at most two positions in a containing progression. It does not improve the general dependence on c_0 in Proposition 7.3.

II.309.2 Torsion-free abelian groups

Theorem II.309.2 (Transfer of the energy classification). *Let G be a torsion-free abelian group and $A \subset G$ have $m \geq 11$ points. The first-four-level classification and Corollary II.309.1 remain*

valid, with models interpreted as

$$b + Sv = \{b + tv : t \in S\}, \quad v \in G \setminus \{0\},$$

where S is one of the integer templates. The third-level equality classification similarly holds for $m \geq 8$.

Proof. The subgroup generated by the finite set A is a finitely generated torsion-free abelian group, hence isomorphic to $\mathbb{Z}^d$ for some finite d . Choose d real numbers linearly independent over $\mathbb{Q}$. Sending a basis to these numbers gives an injective additive map into $\mathbb{R}$. It preserves every additive quadruple in both directions, so the energy and cardinality of the image equal those of A .

Apply the real classification to the image. Every displayed integer template has two points whose parameters differ by one. Its real common step is therefore an actual difference of two image points, and lifts to a group element v . A model point with parameter zero lifts to b . Injectivity then pulls back every relation $x = b + tv$, giving the whole model in G . The progression and difference-set conclusions follow from these same relations.

Only the finitely generated subgroup containing A is embedded. No embedding of an arbitrary ambient torsion-free group into $\mathbb{R}$ is assumed. $\square$

II.309.3 A graph-rigidity corollary for the inverse step

Corollary II.309.3 (Affine graphs in the four-level window). *Let $S \subset \mathbb{Z}$ have $m \geq 11$ elements, let $\phi : S \rightarrow \mathbb{Z}^d$, and let*

$$\Gamma = \{(x, \phi(x)) : x \in S\} \subset \mathbb{Z}^{d+1}.$$

If $E(\Gamma) \geq M_m - 8m + 20$, then Γ is a rank-one arithmetic template from Table II.308.1. In particular, there exist $x_0 \in \mathbb{Z}$, $y_0 \in \mathbb{Z}^d$, $u \in \mathbb{Z} \setminus \{0\}$, $v \in \mathbb{Z}^d$, and a listed parameter set I such that

$$S = \{x_0 + tu : t \in I\}, \quad \phi(x_0 + tu) = y_0 + tv \quad (t \in I).$$

Thus ϕ is the restriction of a rational affine map on its entire domain S .

Proof. Apply Theorem II.309.2 to Γ . Write the direction of the resulting one-dimensional model as (u, v) . If $u = 0$, two different graph points in that model have the same first coordinate, contradicting that Γ is a graph. Thus $u \neq 0$, and the stated formulas follow. Eliminating t gives $\phi(x) = y_0 + (x - x_0)v/u$ on S . $\square$

This corollary concerns a torsion-free graph with an exceptionally high energy. In Gowers's proof the selected graph often lives in a cyclic ambient group, and the available energy need not be close to this endpoint. Both changes of hypothesis must be justified before using the corollary in that setting. In particular, it supplies no missing rectification theorem by itself.

II.309.4 A safe cyclic lifting interface

Proposition II.309.4 (No-wrap transfer). *Let q be a positive integer and let $A \subset \mathbb{Z}/q\mathbb{Z}$ have a set of distinct integer representatives $\tilde{A} \subset [0, L]_{\mathbb{Z}}$ with $q > 2L$. Then*

$$E_{\mathbb{Z}/q\mathbb{Z}}(A) = E_{\mathbb{Z}}(\tilde{A}).$$

Consequently, every real or integer classification above applies to $\tilde{A}$ when its corresponding hypotheses hold, and reduces to a progression-with-holes description of A modulo q .

Proof. For four representatives, the difference between the two pair sums lies in $[-2L, 2L]$. A multiple of q in that interval is zero since $q > 2L$. Thus an additive congruence holds exactly when the integer additive equality holds. $\square$

The strict inequality $q > 2L$ is essential to this simple argument; $q > L$ alone does not prevent wraparound between pair sums. A finite subgroup also shows why no torsion-free endpoint statement can be transferred without a hypothesis: it has energy m^3 , exceeding M_m for $m > 1$.

II.309.5 What would be needed to improve a global progression bound

The local theorem isolates a possible terminal branch in an inverse argument: if a rectified graph is already in the four-level window, its global affine structure is exact. To improve an upper bound for Szemerédi's theorem, one would still need a quantitative mechanism forcing that branch, or showing that the complementary branch yields a more efficient density increment. Neither assertion is proved here. The all- r normal form may be useful in an energy-peeling argument because the template count depends on the local defect rather than on the ambient size; turning this observation into a complete iteration requires additional estimates.

II.310 Sources 68, 62, 61 and 60: Lemma 9.3 on every modulus [write]

Added in the eleventh write, 7 October 2026.

The next twenty sections are source 68's Sections 5–7, source 62's Sections 2–9, source 61's Sections 2–9 and source 60's Section 9, on the restriction step of Gowers's Lemma 9.3 (and, for source 61, the Freiman restriction of Lemma 7.5). *Logically they continue source 58's chapter* on Lemma 9.3 on every cyclic group (Section II.276), with sources 17, 29 and 48 before it (Sections II.57, II.119, II.186), and, for source 61, source 59's character separation (Section II.261); they are printed here because no number printed earlier may move. Their introductions are Sections F.73, F.67, F.66 and F.65.

What is proved. *Source 68:* for B in a finite cyclic group, $b = |B|$, $\phi : B \rightarrow H$ into any abelian group with $G_{16}(B, \phi) \geq \alpha b^{15}$ and $b \geq (4/(\alpha\eta))^{53}$, a subset B' with $T_{16}(B') \geq (\alpha\eta/4)^{53}b^{15}$ and $G_{16}(B') \geq (1 - \eta)T_{16}(B')$ (Theorem II.312.2), by source 58's expectation theorem (reproved, Theorem II.311.5) with a degree-twelve rational filter and exact integer certificates (Proposition II.312.1); an asymmetric form $2^{-205/2}\alpha^{105/2}\eta^{103/2}$ above $b \geq 576 \cdot 2^{94}(\alpha\eta)^{-49}$ (Corollary II.312.3); and two-parameter bounds on the principal cost $\mathcal{I}_{m,a}$, with $\mathcal{I}_{m,a}/m \rightarrow 1 + 1/(2c)$ when $a = m^{c+o(1)}$ (Theorem II.313.1, Corollary II.313.6). *Source 62:* for $B \subseteq \mathbb{Z}/N\mathbb{Z}$ of size βN and $\phi : B \rightarrow \mathbb{Z}/M\mathbb{Z}$, unit-dilation averaging gives restriction with no relation between N and M (Theorem II.314.1); explicitly, retention $2^{-245}\alpha^{31}\eta^{30}|B|^{15}$ above $N \geq 2^{253}(\alpha\eta)^{-31}\beta^{-15}$ (Corollary II.319.2), and $\alpha^{16}\eta^{15}$ divided by a double logarithm in the refined form (Corollary II.314.2), which is necessary for its detector (Theorem II.320.2); the exact weighted degeneracy constant with limit $\zeta(m-2)/\zeta(m-1)$, not a uniform bound (Theorem II.317.1, Proposition II.317.3). *Source 61:* on arbitrary finite abelian groups, $|\Gamma - \Gamma| \leq K|\Gamma|$ partitions the whole domain into at most $5mK^{(2m+1)\log_2(5m)}$ cells on which ϕ is an exact Freiman homomorphism of order m (Theorem F.66.1); a threshold-free exact order-eight restriction with $2^{-2193}\alpha^{911}|B|$ points and $2^{-35277}\alpha^{14655}|B|^{15}$ tuples (Theorem F.66.2); and $\lfloor \log_2 m \rfloor + 1 \leq b_m \leq \log_2(5m)$ for the universal avoidance-colouring exponent

(Theorem F.66.3). *Source 60*: a translated Fejér selector with an all-group certificate (Theorem II.330.1, Corollary II.330.3) gives, on prime moduli, source 17's retention $\alpha^m \eta^{m-1} / (2(m+2)^m)$ under $b \geq 2(1 + \binom{m}{2})L^{m-1}$, $L = \lceil (m+1)/(\alpha\eta) \rceil$ (Corollary II.330.5).

Which is best in which setting (checked at the write; the inequalities in exact arithmetic). (i) *Prime moduli*, $\mathbb{F}_p \rightarrow \mathbb{F}_p$: source 60's corollary keeps the powers $\alpha^{16}\eta^{15}$ of source 17 at $m = 16$ with the threshold $242L^{15}$, one power of L below the intake's marked composition $2(1 + \binom{m}{2})L^m$ (item II.6); on fixed fields source 29 remains better. (ii) *Cyclic domain and cyclic target*, $\mathbb{Z}/N\mathbb{Z} \rightarrow \mathbb{Z}/M\mathbb{Z}$: source 62's powers $\alpha^{31}\eta^{30}$ are the best explicit ones; against source 68's $(\alpha\eta/4)^{53}$ its retention is larger exactly when $2^{-245}\alpha^{31}\eta^{30} > 2^{-106}\alpha^{53}\eta^{53}(1-\eta)$, that is $\alpha^{22}\eta^{23}(1-\eta) < 2^{-139}$ (counting respected tuples on both sides), while its threshold $N \geq 2^{253}(\alpha\eta)^{-31}\beta^{-15}$ grows like β^{-15} where source 68's, $N \geq \beta^{-1}(4/(\alpha\eta))^{53}$, grows like β^{-1} ; neither dominates. (iii) *Cyclic domain, arbitrary abelian target*: source 68's exponent 53 is the best (source 58: 54); source 62 does not apply (its question on noncyclic targets, Section V.235). (iv) *Arbitrary finite abelian domain and target*: only source 61 applies; its restriction is exact and threshold-free, at the cost of the powers α^{911} and α^{14655} . All of them improve the printed exponent 2^{19} . *Against the catalogue* (`lemma_9_3`, `Sections08_09.lean`, line 77, all $N \geq N_0$, maps $\mathbb{Z}/N\mathbb{Z} \rightarrow \mathbb{Z}/N\mathbb{Z}$, retention $(\alpha\eta/4)^{2^{19}}\beta^{15}N^{15}$): source 68's theorem implies it with $N_0 = \lceil \beta^{-1}(4/(\alpha\eta))^{53} \rceil$, source 62's with $N_0 = \lceil 2^{253}(\alpha\eta)^{-31}\beta^{-15} \rceil$ ($M = N$), and source 61's with $N_0 = 1$; source 61's Corollary II.328.2 also implies `corollary_9_4` with $N_0 = 1$ (notes at the corollaries). Source 60's is for prime moduli only. Both catalogue statements are already exact (`lemma_9_3_holds`, `corollary_9_4_holds`, `Proofs09Restriction.lean`, lines 376 and 416), and none of these theorems is formalized.

Credits. Source 68 reproves and credits source 58's expectation theorem; source 62's pin precedes source 58's placement, so their overlap is not a lapse; source 61 finds source 59's exponent $(2m+1)\log_2(5m)$ independently, and source 59's bound is sharper by the factor $5m$ and in its defect-sensitive form (note after Theorem F.66.1); source 61's Plünnecke argument is Petridis's and its energy-to-growth step a conservative Balog–Szemerédi–Gowers argument, credited; source 60's selector uses source 17's Fejér kernel and source 04's weighted rank-two count, credited.

Relation to the formal statements (checked at HEAD fb9602e55). `lemma_9_3` and `corollary_9_4` (`Sections08_09.lean`, lines 77 and 91) and `lemma_7_5`, `corollary_7_6` (`Sections06_07.lean`, lines 183 and 197) have exact companions (`lemma_9_3_holds`, `corollary_9_4_holds` in `Proofs09Restriction.lean`, lines 376 and 416; `lemma_7_5_holds`, `Proofs07FreimanRestriction.lean`, line 267; `corollary_7_6_holds`, `Proofs07AdditiveRestriction.lean`, line 100); the tuple counts are `additiveTupleCount`, `phiAdditiveTupleCount` and `GammaHomOfOrder` (`Definitions.lean`, lines 322, 327 and 334). Nothing in this chapter is formalized.

II.311 A restriction estimate with decaying exceptional terms

Source 68, delivered Section 5.

The new finite filter in Section II.312 is inserted into a restriction theorem already established in source 58 [281]. We give its proof here so that the numerical improvement can be checked without importing an unspecified selection argument. The underlying random-character restriction is Gowers' method [1, Lemma 9.3]; the weighted nonprincipal estimate and the decaying correction for repeated points are credited to source 58.

II.311.1 Ordered counts and admissible filters

For this section fix an even integer $m = 2r \geq 4$ and write

$$\varepsilon = (\underbrace{1, \dots, 1}_r, \underbrace{-1, \dots, -1}_r).$$

Initially G, H are finite abelian groups, $B \subseteq G$ is nonempty, $b = |B|$, and $\phi : B \rightarrow H$. Define

$$T_m(B) = |\{\mathbf{x} \in B^m : \sum_i \varepsilon_i x_i = 0\}|, \quad (\text{II.311.1})$$

$$G_m(B, \phi) = |\{\mathbf{x} \in B^m : \sum_i \varepsilon_i x_i = 0, \sum_i \varepsilon_i \phi(x_i) = 0\}|, \quad (\text{II.311.2})$$

$$U_m(B, \phi) = T_m(B) - G_m(B, \phi). \quad (\text{II.311.3})$$

Every count is of ordered tuples, with repetitions allowed. An additive tuple is *respected* when its target discrepancy $\sum_i \varepsilon_i \phi(x_i)$ is zero, and *bad* otherwise.

Lemma II.311.1 (Free variables and repeated entries). *One has $T_m(B) \leq b^{m-1}$. The number of additive tuples having a repeated entry is at most $\binom{m}{2} b^{m-2}$. Every bad tuple has at least two distinct entries.*

Proof. The first $m-1$ entries determine the last, whose coefficient is ± 1 . To count repetitions, fix positions $i < j$ with $x_i = x_j$ and choose a position h outside that pair. Choose the entries in the $m-2$ positions other than j, h . The equality determines x_j , and the additive relation determines x_h . Thus there are at most b^{m-2} tuples for each pair; a union bound proves the assertion. If every entry is the same point, both balanced target sums equal $r\phi(x)$, so that tuple is respected. $\square$

Put $e(t) = \exp(2\pi it)$ and identify $\mathbb{T}$ with $\mathbb{R}/\mathbb{Z}$, equipped with Haar measure of total mass one. A *decreasing admissible filter of degree D* is a real trigonometric polynomial

$$K(t) = \sum_{n=-D}^D c_n e(nt), \quad 0 \leq K \leq 1, \quad K(0) = 1, \quad (\text{II.311.4})$$

whose coefficients satisfy

$$c_{-n} = c_n \geq 0, \quad c_0 \geq c_1 \geq \dots \geq c_D, \quad c_1 > 0. \quad (\text{II.311.5})$$

Coefficients outside the indicated range are zero. Define

$$P = \sum_n c_n^m, \quad P_a = \sum_{a|n} c_n^m \quad (a \geq 2), \quad Q = P_2, \quad R_j = \int_{\mathbb{T}} K(t)^j dt. \quad (\text{II.311.6})$$

Here a is an integer order, whereas the subscript j in R_j denotes a power of the filter. In particular,

$$0 < Q < P < 1, \quad P \leq R_m \leq R_2 \leq 1. \quad (\text{II.311.7})$$

Indeed, $\sum c_n = 1$, at least c_{-1}, c_0, c_1 are positive, and the sum defining Q omits $c_{-1}^m + c_1^m$. The expansion of the constant coefficient of K^m contains all vectors $(t\varepsilon_1, \dots, t\varepsilon_m)$, of total weight P , proving $P \leq R_m$. The remaining comparisons follow from $0 \leq K \leq 1$.

Lemma II.311.2 (Target orders). *If $a \geq a_0 \geq 2$, then $P_a \leq P_{a_0}$. If an element $z \in H$ has order a , then*

$$\sum_{n \in \mathbb{Z}} c_n^m \mathbf{1}_{nz=0} = P_a \leq Q.$$

Proof. Evenness and monotonicity give $P_a = c_0^m + 2 \sum_{j \geq 1} c_{aj}^m \leq c_0^m + 2 \sum_{j \geq 1} c_{a_0 j}^m$. The exact formula follows from $nz = 0$ if and only if $a \mid n$. $\square$

II.311.2 A finite random selector

For $x \in G$, normalized character orthogonality is

$$\mathbb{E}_{\chi \in \widehat{G}} e(\chi(x)) = \mathbf{1}_{x=0}. \quad (\text{II.311.8})$$

For $x \neq 0$, the decomposition of a finite abelian group into cyclic groups supplies a character that is nontrivial on x . Translating the character sum by it proves that the sum vanishes.

Choose integers $k \geq 1$ and $M > mD$. Independently for each $\ell = 1, \dots, k$, choose uniform characters $\chi_\ell \in \widehat{G}$, $\psi_\ell \in \widehat{H}$, and a uniform shift $u_\ell \in \{0, 1/M, \dots, (M-1)/M\}$. Conditional on these parameters, select each distinct point $x \in B$ independently with probability

$$p(x) = \prod_{\ell=1}^k K(\chi_\ell(x) + \psi_\ell(\phi(x)) + u_\ell). \quad (\text{II.311.9})$$

This defines a random subset B' . The parameter choices and the set of possible subsets are finite; the probabilities need not be rational.

For one layer, expansion and orthogonality give the identity

$$\begin{aligned} & \mathbb{E}_{\chi, \psi, u} \prod_{i=1}^m K(\chi(x_i) + \psi(\phi(x_i)) + u) \\ &= \sum_{\mathbf{n} \in [-D, D]^m} \left(\prod_i c_{n_i} \right) \mathbf{1}_{\sum_i n_i = 0} \mathbf{1}_{\sum_i n_i x_i = 0} \mathbf{1}_{\sum_i n_i \phi(x_i) = 0}. \end{aligned} \quad (\text{II.311.10})$$

The first equality tested on the right is in $\mathbb{Z}$. The grid average initially imposes divisibility by M , and $|\sum_i n_i| \leq mD < M$ makes divisibility equivalent to zero. The other equalities are in G and H .

Call a coefficient vector *principal* if $\mathbf{n} = t\epsilon$ for an integer t . For a respected additive tuple, every principal vector passes all three tests, with total weight P . For a bad additive tuple of discrepancy order a , its principal weight is P_a . All coefficients in the expansion are nonnegative.

For tuples with distinct entries, the product in (II.311.10), multiplied over layers, is the actual conditional containment probability. For repeated entries it is a *formal product over occurrences*. If x occurs μ_x times, the formal factor is $p(x)^{\mu_x}$ and the actual factor is $p(x)$. Consequently actual containment is at least formal containment. It follows that

$$\mathbb{E} G_m(B', \phi) \geq P^k G_m(B, \phi), \quad (\text{II.311.11})$$

including all repetitions. This direction of comparison cannot be used to upper-bound bad repeated tuples; those require the next estimates.

II.311.3 Nonprincipal relations and collision errors

For $d \neq 0$ put $G[d] = \{x \in G : dx = 0\}$, and define

$$\tau_G(L) = \max_{1 \leq d \leq L} |G[d]|. \quad (\text{II.311.12})$$

For cyclic groups, including the trivial group,

$$|(\mathbb{Z}/N\mathbb{Z})[d]| = \gcd(N, d), \quad \tau_{\mathbb{Z}/N\mathbb{Z}}(2D) \leq 2D. \quad (\text{II.311.13})$$

Lemma II.311.3 (Completion with two relations). *For a nonprincipal $\mathbf{n} \in [-D, D]^m \cap \mathbb{Z}^m$, at most $\tau_G(2D)b^{m-2}$ tuples in B^m satisfy both $\sum_i \varepsilon_i x_i = 0$ and $\sum_i n_i x_i = 0$.*

Proof. There are positions i, j with $d = \varepsilon_i n_j - \varepsilon_j n_i \neq 0$. Otherwise all integers n_i/ε_i would be equal, making the vector principal. Moreover $|d| \leq 2D$. Fix the other $m - 2$ entries. The first relation, whose relevant coefficient is ± 1 , expresses one remaining entry in terms of the other. Substitution into the second gives an equation $dx = z$, up to an immaterial sign. Every nonempty multiplication fibre is a coset of $G[d]$. Restriction of both entries to B only reduces the number of completions. $\square$

The total coefficient weight of one layer satisfying just the integer zero-sum test is the constant coefficient of K^m , namely R_m . Its principal part has weight P . Hence the total weight of all k -layer coefficient arrays containing a nonprincipal layer is $R_m^k - P^k$. For each such array choose one nonprincipal layer and apply Lemma II.311.3, discarding all other conditions. Summing nonnegative weights proves that its total contribution over distinct-entry additive tuples is at most

$$\tau_G(2D)(R_m^k - P^k)b^{m-2}. \quad (\text{II.311.14})$$

Each array is counted once. There is no union bound over the layers and no additional factor k .

Lemma II.311.4 (Containment of a tuple with repeated entries). *A tuple with $d \geq 2$ distinct entries has containment probability at most $R_d^k \leq R_2^k$.*

Proof. Fix the characters in one layer and write $\theta_1, \dots, \theta_d$ for the phases of its distinct points. Hölder's inequality gives

$$\mathbb{E}_u \prod_{j=1}^d K(u + \theta_j) \leq \prod_{j=1}^d (\mathbb{E}_u K(u + \theta_j)^d)^{1/d} = R_d.$$

To justify the equality, each translated polynomial has degree at most $dD \leq mD < M$; its grid average therefore equals its constant coefficient, which is R_d . Average over the characters and multiply over the independent layers. Finally $K^d \leq K^2$ for $d \geq 2$. $\square$

By Lemma II.311.1, the expected number of retained bad tuples with repeated entries is at most

$$\binom{m}{2} R_2^k b^{m-2}. \quad (\text{II.311.15})$$

Theorem II.311.5 (Repeated-filter estimate; source 58). *For the selector (II.311.9),*

$$\mathbb{E} G_m(B', \phi) \geq P^k G_m(B, \phi), \quad (\text{II.311.16})$$

$$\mathbb{E} U_m(B', \phi) \leq Q^k U_m(B, \phi) + A_k b^{m-2}, \quad (\text{II.311.17})$$

where

$$A_k = \tau_G(2D)(R_m^k - P^k) + \binom{m}{2} R_2^k \leq C_G R_2^k, \quad C_G = \tau_G(2D) + \binom{m}{2}. \quad (\text{II.311.18})$$

Proof. The good bound is (II.311.11). For a bad distinct-entry tuple, the all-principal part is at most Q^k by Lemma II.311.2. Sum that part over the bad tuples, add (II.311.14), and then add (II.311.15). The final inequality uses $R_m \leq R_2$ and drops the nonpositive term $-\tau_G(2D)P^k$. $\square$

II.311.4 One subset with both quality and retained mass

Theorem II.311.6 (Spectral-gap restriction; source 58). *Suppose $0 < \alpha, \eta \leq 1$ and $G_m(B, \phi) \geq \alpha b^{m-1}$. Put*

$$t = \alpha\eta, \quad \mathcal{R} = P/Q > 1, \quad k = \left\lceil \frac{\log(2/t)}{\log \mathcal{R}} \right\rceil. \quad (\text{II.311.19})$$

If

$$b \geq \frac{4A_k}{tP^k}, \quad (\text{II.311.20})$$

then some $B' \subseteq B$ satisfies

$$G_m(B', \phi) \geq (1 - \eta)T_m(B'), \quad T_m(B') \geq \frac{\alpha P^k}{4} b^{m-1}. \quad (\text{II.311.21})$$

A sufficient replacement for (II.311.20) is $b \geq 4C_G t^{-1} (R_2/P)^k$.

Proof. If $\eta = 1$, take $B' = B$ and use $T_m(B) \geq G_m(B, \phi)$. Assume $\eta < 1$. For each outcome consider the single random variable

$$F = G_m(B', \phi) - (1 - \eta)T_m(B') = \eta G_m(B', \phi) - (1 - \eta)U_m(B', \phi).$$

The choice of k implies $Q^k \leq (t/2)P^k$. Since $U_m(B, \phi) \leq b^{m-1}$, Theorem II.311.5 gives

$$\begin{aligned} \mathbb{E}F &\geq tP^k b^{m-1} - (1 - \eta)Q^k b^{m-1} - (1 - \eta)A_k b^{m-2} \\ &\geq \frac{t}{2}P^k b^{m-1} - A_k b^{m-2} \geq \frac{t}{4}P^k b^{m-1} > 0. \end{aligned}$$

Choose an outcome with $F \geq \mathbb{E}F$. Its positivity proves the quality claim, and $F = \eta T_m(B') - U_m(B', \phi) \leq \eta T_m(B')$ proves the retained-mass claim. The simpler threshold follows from (II.311.18). $\square$

Define the two real exponents

$$a_K = \frac{-\log P}{\log \mathcal{R}}, \quad z_K = \frac{\log(R_2/P)}{\log \mathcal{R}}. \quad (\text{II.311.22})$$

Here $a_K > 0$ and $z_K \geq 0$. Since $k < \log(2/t)/\log \mathcal{R} + 1$,

$$P^k \geq P(t/2)^{a_K}, \quad (R_2/P)^k \leq (R_2/P)(2/t)^{z_K}. \quad (\text{II.311.23})$$

Thus a sufficient threshold and the resulting retention are

$$b \geq 4C_G(R_2/P) 2^{z_K} (\alpha\eta)^{-(z_K+1)}, \quad (\text{II.311.24})$$

$$T_m(B') \geq \frac{P}{4} 2^{a_K} \alpha^{a_K+1} \eta^{a_K} b^{m-1}. \quad (\text{II.311.25})$$

These are the precise interfaces used below. The powers of α and η need not be rounded to the same integer.

II.311.5 Arbitrary abelian target groups

Lemma II.311.7 (A finite quotient preserving the tuple relations). *For finite B , any map $\phi : B \rightarrow H$ to an abelian group, and a fixed m , there is a finite abelian group H_0 and a map $\phi_0 : B \rightarrow H_0$ preserving precisely which balanced m -term target discrepancies vanish. The same equivalence holds on every subset of B .*

Proof. The subgroup generated by $\phi(B)$ is finitely generated, hence is isomorphic to $\mathbb{Z}^s \oplus F$ for a finite abelian group F . List all balanced m -term discrepancies of points of B . Choose an integer $L \geq 2$ strictly larger than the absolute values of all coordinates of their free parts. Reduce the free factor modulo L and retain F . A nonzero free coordinate remains nonzero modulo L ; a nonzero discrepancy with zero free part retains its nonzero torsion coordinate. Zero discrepancies remain zero. Every tuple in a subset was already in the original finite list. $\square$

Theorem II.311.6 and (II.311.24)– (II.311.25) therefore hold for arbitrary abelian targets, with the same constants. Only the domain is required to be finite. The general-domain torsion factor remains present: for $G = (\mathbb{Z}/2\mathbb{Z})^s$, already $|G[2]| = |G|$. Its bound by $2D$ is a specific advantage of cyclic domains, rather than a property of all finite abelian groups.

II.312 A degree-twelve filter and the exponent fifty-three

Source 68, delivered Section 6.

The restriction proof is now fixed. Its remaining inputs are the coefficients of a filter and exact inequalities between their moments. Source 58 uses $\cos^{50}(\pi t)$, of degree 25, to obtain the common exponent 54 in its size threshold and retained tuple fraction [281]. The following rational filter reduces the degree to 12 and the common exponent to 53. It also gives a principal cost strictly below $103/2$, compared with approximately 52.2064 for that specific powered cosine. No optimum over all filters is asserted.

II.312.1 An admissible polynomial with rational coefficients

Extend the following sequence evenly and by zero outside $[-6, 6]$:

$$(a_0, a_1, \dots, a_6) = (1000, 942, 782, 558, 320, 124, 14). \quad (\text{II.312.1})$$

Set

$$S = \sum_{j=-6}^6 a_j = 6480, \quad H(t) = \sum_{j=-6}^6 a_j e(jt), \quad K(t) = \frac{H(t)^2}{S^2}. \quad (\text{II.312.2})$$

Since the sequence is real and even, H is real. Thus $K \geq 0$; the triangle inequality gives $|H(t)| \leq S$, so $K \leq 1$; and $H(0) = S$ gives $K(0) = 1$. Write

$$d = S^2 = 41990400, \quad b_n = \sum_j a_j a_{n-j}, \quad c_n = b_n/d. \quad (\text{II.312.3})$$

The c_n are the Fourier coefficients of K and are nonnegative and even. Their positive-side data are small enough to print in full.

n	0	1	2	3	4	5
b_n	4856448	4669952	4150460	3401728	2558308	1749968
n	6	7	8	9	10	11
b_n	1073460	577432	262680	94984	24336	3472

The final coefficient is $b_{12} = 196$. In particular, $b_0 > b_1 > \cdots > b_{12} > 0$ and $b_1 > 0$. Also $\sum_{n=-12}^{12} b_n = S^2$. This proves all the conditions of (II.311.4)–(II.311.5), with degree exactly $D = 12$.

II.312.2 Exact integer certificates

For the rest of this section $m = 16$. Introduce the integers

$$A = \sum_{n=-12}^{12} b_n^{16}, \quad B_{\text{ev}} = \sum_{\substack{-12 \leq n \leq 12 \\ 2|n}} b_n^{16}, \quad V = \sum_{n=-12}^{12} b_n^2. \quad (\text{II.312.4})$$

Then

$$P = \frac{A}{d^{16}}, \quad Q = \frac{B_{\text{ev}}}{d^{16}}, \quad \mathcal{R} = \frac{A}{B_{\text{ev}}}, \quad R_2 = \frac{V}{d^2}. \quad (\text{II.312.5})$$

The last identity is the constant coefficient of K^2 , since $c_{-n} = c_n$. In particular $V = 147141531710208$.

Proposition II.312.1 (Finite certificates). *The integers in (II.312.4) satisfy all the inequalities in the following table.*

<i>Exact inequality</i>	<i>Consequence</i>
$2^{49}A > d^{16}$	$P > 2^{-49}$
$250A > 481B_{\text{ev}}$	$\mathcal{R} > 481/250$
$481^{52} > 2^{49}250^{52}$	$\mathcal{R}^{52} > 2^{49}$
$481^{49} > 2^{46}250^{49}$	$\mathcal{R}^{49} > 2^{46}$
$8V < d^2$	$R_2 < 2^{-3}$
$A^{105} > d^{32}B_{\text{ev}}^{103}$	$a_K < 103/2$
$A^{49} > B_{\text{ev}}^{48}Vd^{14}$	$z_K < 48$

Consequently $a_K < 52$ and $z_K < 49$ also hold.

Proof. The inequalities are finite integer arithmetic with the explicitly listed coefficients. To make the certificate independently executable, the complete essential computation is printed below; exponentiation uses integers of arbitrary precision.

```
seed = [1000, 942, 782, 558, 320, 124, 14]
a = {j: seed[abs(j)] for j in range(-6, 7)}
d = sum(a.values()) ** 2
b = {n: sum(a[j] * a.get(n-j, 0) for j in a)
     for n in range(-12, 13)}
A = sum(v ** 16 for v in b.values())
B = sum(v ** 16 for n, v in b.items() if n % 2 == 0)
V = sum(v * v for v in b.values())
```

```

checks = [
    d == 41990400,
    all(b[n] > b[n+1] for n in range(12)),
    2**49 * A > d**16,
    250 * A > 481 * B,
    481**52 > 2**49 * 250**52,
    481**49 > 2**46 * 250**49,
    8 * V < d**2,
    A**105 > d**32 * B**103,
    A**49 > B**48 * V * d**14,
]
if not all(checks):
    raise ArithmeticError("Certificate failed")

```

The accompanying verifier executes these comparisons and additional normalization checks. No floating-point inequality enters the certificate.

For the last two consequences, the penultimate inequality is precisely $P^2\mathcal{R}^{103} > 1$, implying $-\log P/\log \mathcal{R} < 103/2$. The final inequality is precisely $\mathcal{R}^{48}P/R_2 > 1$, implying $\log(R_2/P)/\log \mathcal{R} < 48$. These arguments only take logarithms of strictly positive numbers and use $\mathcal{R} > 1$. $\square$

For orientation, the values are

$$\begin{aligned}
 P &\approx 2.2934630227959095 \cdot 10^{-15}, & \mathcal{R} &\approx 1.925599333199833, \\
 a_K &\approx 51.44504972230757, & z_K &\approx 47.654839268529635.
 \end{aligned} \tag{II.312.6}$$

These decimals describe the certified filter. The theorem below rests on Proposition II.312.1.

II.312.3 The resulting restriction theorem

Theorem II.312.2 (A common exponent of fifty-three). *Let G be any finite cyclic group, let H be any abelian group, and let $\phi : B \rightarrow H$ for a nonempty $B \subseteq G$ of cardinality b . Suppose $0 < \alpha, \eta \leq 1$ and*

$$G_{16}(B, \phi) \geq \alpha b^{15}.$$

If

$$b \geq \left(\frac{4}{\alpha\eta}\right)^{53}, \tag{II.312.7}$$

then a subset $B' \subseteq B$ satisfies

$$\begin{aligned}
 G_{16}(B', \phi) &\geq (1 - \eta)T_{16}(B'), \\
 T_{16}(B') &\geq \left(\frac{\alpha\eta}{4}\right)^{53} b^{15}, & |B'| &\geq \left(\frac{\alpha\eta}{4}\right)^{53/15} b.
 \end{aligned} \tag{II.312.8}$$

The counts include repeated entries and require no restriction on the prime factors of $|G|$.

Proof. First assume H finite and put $t = \alpha\eta$. The filter has $D = 12$, hence (II.311.13) gives $C_G \leq 24 + \binom{16}{2} = 144$. The first five certificates imply

$$P \geq 2^{-49}, \quad R_2/P \leq 2^{46}, \quad a_K \leq 52, \quad z_K \leq 49.$$

Thus the sufficient threshold (II.311.24) is at most

$$576 2^{95} t^{-50} \leq 2^{106} t^{-53} = (4/t)^{53}.$$

Here $576 \leq 2^{11}$ and $0 < t \leq 1$. The retained proportion in (II.311.25) is at least

$$\frac{2^{-49}}{4 2^{52}} \alpha^{53} \eta^{52} = 2^{-103} \alpha^{53} \eta^{52} \geq 2^{-106} \alpha^{53} \eta^{53}.$$

The inequalities in the powers use $0 < \alpha, \eta \leq 1$. Theorem II.311.6 supplies the same subset for quality and mass. Lemma II.311.1 gives $T_{16}(B') \leq |B'|^{15}$, yielding the size conclusion. For arbitrary H , apply Lemma II.311.7 and transfer the selected subset back to the original map. $\square$

[write] **7 October 2026; sources 60–69.** *Against the formal statement and the other all-moduli theorems* (checked at the write). With $b = \beta N$ and $N_0 = \lceil \beta^{-1} (4/(\alpha\eta))^{53} \rceil$ this theorem gives `lemma_9_3` (Sections08_09.lean, line 77), since $0 < \alpha\eta/4 < 1$ and $53 < 2^{19}$, on paper; the catalogue statement is already exact and this theorem is not formalized. It improves source 58's exponent 54 (Theorem A) in the same setting; for cyclic targets source 62's Corollary II.319.2 retains more when $\alpha^{22} \eta^{23} (1 - \eta) < 2^{-139}$ (Section II.310).

Corollary II.312.3 (A stronger asymmetric parameter bound). *Under the same count hypothesis, the sufficient size condition*

$$b \geq 576 2^{94} (\alpha\eta)^{-49} \tag{II.312.9}$$

guarantees a subset with respected proportion at least $1 - \eta$ and

$$T_{16}(B') \geq 2^{-205/2} \alpha^{105/2} \eta^{103/2} b^{15}. \tag{II.312.10}$$

Proof. Use $z_K \leq 48$ in (II.311.24), together with $C_G \leq 144$ and $R_2/P \leq 2^{46}$. Use $a_K \leq 103/2$ and $P \geq 2^{-49}$ in (II.311.25). The exponent of two in the denominator is $49 + 2 + 103/2 = 205/2$. $\square$

The comparison with source 58 is on the same cyclic-domain, arbitrary-target interface and with the same tuple conventions: the common threshold $(4/(\alpha\eta))^{54}$ and fraction $(\alpha\eta/4)^{54}$ are both improved by replacing 54 with 53. The asymmetric corollary retains more of the certified information. These are local restriction bounds. No propagation through the full inverse argument or density-increment iteration is supplied here.

II.313 Uniform principal costs as arity and torsion vary

Source 68, delivered Section 7.

The preceding restriction uses discrepancy order two as its worst case. When all nonzero target discrepancies have order at least a , the principal bad weight can instead be P_a . Source 58 proves, for each fixed $a \geq 2$, the asymptotic

$$\mathcal{I}_{m,a} = \frac{m \log m}{2 \log a} + O_a(m \log \log m) \quad (m \rightarrow \infty)$$

for the principal cost defined below [281]. We obtain a uniform estimate, allow a to grow with m , and improve the fixed- a error. An additional term $m - 1$ becomes a leading term when a is a power of m .

II.313.1 The extremal problem and its two-parameter answer

In this section $m \geq 2$ is any integer. Evenness of m is needed for its application to a balanced relation, but is unnecessary for the extremal problem. Let $\mathcal{K}$ be the class of all finite trigonometric polynomials

$$K(x) = \sum_n c_n e(nx), \quad 0 \leq K \leq 1, \quad K(0) = 1, \quad c_n = c_{-n} \geq 0. \quad (\text{II.313.1})$$

Monotonicity of the coefficients is not required in this class. Put

$$P_m(K) = \sum_n c_n^m, \quad Q_{m,a}(K) = \sum_{a|n} c_n^m \quad (a \geq 2), \quad (\text{II.313.2})$$

and define

$$\mathcal{A}_{m,a}(K) = \frac{-\log P_m(K)}{\log(P_m(K)/Q_{m,a}(K))}, \quad \mathcal{I}_{m,a} = \inf_{K \in \mathcal{K}} \mathcal{A}_{m,a}(K). \quad (\text{II.313.3})$$

Zero gain $P_m = Q_{m,a}$ is assigned cost $+\infty$. All logarithms are natural.

Theorem II.313.1 (A uniform arity–torsion estimate). *For every pair of integers $m, a \geq 2$,*

$$\mathcal{I}_{m,a} \geq m - 1, \quad (\text{II.313.4})$$

$$\mathcal{I}_{m,a} \geq m - 1 + \frac{m(\log m - 20)}{2 \log a}, \quad (\text{II.313.5})$$

$$\mathcal{I}_{m,a} \leq m - 1 + \frac{m}{2 \log a} \left[\log m + \log \log \left(e + \frac{\log m}{\log a} \right) + 4 \right]. \quad (\text{II.313.6})$$

The constants are absolute and the inequalities hold for all the stated finite parameters. No optimality is claimed for the constants 20 and 4.

The lower estimate imports the classical circle Riesz–Sobolev inequality [242, Theorem 1.1 and (1.8)]. The remaining steps, including its precise convolution consequence and the finite approximants in the upper construction, are proved below.

II.313.2 Elementary bounds and a small-gain estimate

Write $s = c_0 = \int_{\mathbb{T}} K$. Since K is continuous, nonnegative, and $K(0) = 1$, one has $s > 0$. If $s = 1$, then $K \equiv 1$ and its cost is infinite. Otherwise $0 < s < 1$. The integral formula for a coefficient and the normalization at zero give

$$0 \leq c_n \leq s, \quad \sum_n c_n = 1, \quad P_m \leq s^{m-1}, \quad Q_{m,a} \geq s^m. \quad (\text{II.313.7})$$

Let $P = P_m$, $Q = Q_{m,a}$, and $R = P/Q$ when the parameters are fixed.

Lemma II.313.2 (Gain ceilings). *For every admissible filter,*

$$1 \leq R \leq a, \quad R \leq s^{-1}.$$

Every positive-gain filter has $\mathcal{A}_{m,a}(K) \geq m - 1$.

Proof. Use normalized convolution $f * g(x) = \int_{\mathbb{T}} f(x-y)g(y) dy$. The function $F = K^{*m}$ is nonnegative and has Fourier coefficients c_n^m . Therefore

$$Q = \frac{1}{a} \sum_{j=0}^{a-1} F(j/a) \geq \frac{F(0)}{a} = \frac{P}{a}.$$

The lower bound $R \geq 1$ follows by deleting nonnegative terms. Equation (II.313.7) gives $R \leq 1/s$. For $R > 1$, $-\log P \geq (m-1)(-\log s)$ and $\log R \leq -\log s$, proving the cost estimate. $\square$

The next elementary estimate is used in source 58. We retain it because a bound for gains bounded away from one is not by itself sufficient to control the infimum in (II.313.3).

Lemma II.313.3 (A coefficient estimate and the small-gain alternative). *For $n \neq 0$,*

$$c_n \leq s \frac{\sin(\pi s)}{\pi s} \leq s \exp(-2s^2/3). \quad (\text{II.313.8})$$

Consequently, for $R > 1$,

$$R - 1 \leq \frac{1-s}{s} \exp\left(-\frac{2(m-1)s^2}{3}\right), \quad (\text{II.313.9})$$

$$\mathcal{A}_{m,a}(K) \geq (m-1)s \exp\left(\frac{2(m-1)s^2}{3}\right). \quad (\text{II.313.10})$$

Proof. The distribution of $\cos(2\pi nx)$ under Haar measure is that of $\cos(2\pi x)$. Let E be its top level set of measure s . The pointwise inequality

$$(\mathbf{1}_E - K)(\cos(2\pi nx) - \cos(\pi s)) \geq 0$$

and $\int(\mathbf{1}_E - K) = 0$ imply $c_n \leq \int_E \cos(2\pi nx) dx = \sin(\pi s)/\pi$. For $0 \leq v \leq 1$, concavity of sine gives $\sin(\pi v/2) \geq v$, hence $\cos(\pi v) \leq 1 - 2v^2$. Integrating with $v = su$ for $0 \leq u \leq 1$ gives $\sin(\pi s)/(\pi s) \leq 1 - 2s^2/3 \leq e^{-2s^2/3}$.

The total mass of the nonzero coefficients is $1 - s$. Thus

$$R - 1 \leq \frac{\sum_{n \neq 0} c_n^m}{s^m} \leq \frac{1-s}{s} \exp\left(-\frac{2(m-1)s^2}{3}\right).$$

Finally use $-\log P \geq (m-1)(-\log s)$, $-\log s \geq 1 - s$, and $\log R \leq R - 1$. $\square$

II.313.3 A convolution-concentration bound

For a nonnegative integrable function f on $\mathbb{T}$, let f^* be its symmetric nonincreasing rearrangement: it is even, nonincreasing on $[0, 1/2]$, and equimeasurable with f . For $0 \leq v \leq 1$, let I_v be the centered interval of measure v . The classical inequality we use is

$$\int_{\mathbb{T}} (f * g)h \leq \int_{\mathbb{T}} (f^* * g^*)h^* \quad (f, g, h \geq 0). \quad (\text{II.313.11})$$

This is the circle case of the cited theorem; that source credits Baernstein for this case [242].

Say that $f \preceq F$, with F symmetric nonincreasing and of the same integral as f , if

$$\int_E f \leq \int_{I_{|E|}} F \quad \text{for every measurable } E \subseteq \mathbb{T}. \quad (\text{II.313.12})$$

Equivalently, the centered partial integrals of f^* are bounded by those of F . By the layer-cake representation, this implies $\int f^* q \leq \int F q$ for every nonnegative symmetric nonincreasing q .

Lemma II.313.4 (Convolution preserves this concentration order). *If $f \preceq F$ and $g \preceq G$, then $f * g \preceq F * G$. In particular, for $0 \leq f \leq 1$ with $\int f = s$ and integers $m \geq 2$,*

$$f^{*m}(x) \leq \mathbf{1}_{I_s}^{*m}(0) \quad \text{for every } x \in \mathbb{T}. \quad (\text{II.313.13})$$

Proof. First, convolution of symmetric nonincreasing functions on the circle is symmetric nonincreasing. By layer cake it suffices to check two interval indicators. If their lengths are α, β , put $p = (\alpha + \beta)/2$ and $q = |\alpha - \beta|/2$. Their real-line overlap at separation $x \geq 0$ is $\ell(x) = (p - x)_+ - (q - x)_+$. On the circle, for $0 \leq x \leq 1/2$, the overlap is $\ell(x) + \ell(1 - x)$. The first summand has slope 0 or -1 . The second can have slope $+1$ only when $x > 1 - p$; then $q \leq 1 - p < x < p$, since $p + q \leq 1$, so the first has slope -1 . Their sum is therefore nonincreasing.

For $|E| = v$, (II.313.11) gives

$$\int_E f * g \leq \int_{I_v} f^* * g^* = \int_{\mathbb{T}} f^*(g^* * \mathbf{1}_{I_v}).$$

The parenthesized function is symmetric nonincreasing. Replace f^* by F using (II.313.12) and layer cake. By commutativity, the resulting expression equals $\int g^*(F * \mathbf{1}_{I_v})$. Replace g^* by G in the same way. The result is $\int_{I_v} F * G$, proving the claim.

If $0 \leq f \leq 1$ and $\int f = s$, then $\int_E f \leq \min(|E|, s) = \int_{I_{|E|}} \mathbf{1}_{I_s}$. Thus $f \preceq \mathbf{1}_{I_s}$, and induction gives $f^{*m} \preceq \mathbf{1}_{I_s}^{*m}$. Apply this to intervals centered at x and let their measure decrease to zero. The functions in this final step are continuous for $m \geq 2$: for bounded f, g , translation continuity in L^1 and $\|f * (g(\cdot + h) - g)\|_\infty \leq \|f\|_\infty \|g(\cdot + h) - g\|_1$ give the needed continuity. $\square$

Proposition II.313.5 (Mean-sensitive spectral concentration). *For every $K \in \mathcal{K}$, $s = \int K$, and integer $m \geq 2$,*

$$P_m(K) \leq s^m \left(1 + \frac{4}{s\sqrt{m}} \right). \quad (\text{II.313.14})$$

Proof. Fourier convolution gives $P_m(K) = K^{*m}(0)$. The Fourier coefficients of $\mathbf{1}_{I_s}$ are $s \operatorname{sinc}(\pi n s)$, where $\operatorname{sinc} z = \sin z / z$ and $\operatorname{sinc} 0 = 1$. Their m th powers are absolutely summable for $m \geq 2$. Lemma II.313.4 therefore implies

$$\frac{P_m(K)}{s^m} \leq 1 + 2 \sum_{n \geq 1} |\operatorname{sinc}(\pi n s)|^m.$$

For $0 \leq x \leq 1$, the sine estimate from Lemma II.313.3 gives $|\operatorname{sinc}(\pi x)| \leq e^{-2x^2/3}$; for $x > 1$ use $|\operatorname{sinc}(\pi x)| \leq (\pi x)^{-1}$. The resulting envelope

$$F_m(x) = \begin{cases} e^{-2mx^2/3}, & 0 \leq x \leq 1, \\ (\pi x)^{-m}, & x > 1 \end{cases}$$

is nonincreasing, including its downward jump at one. Thus

$$\begin{aligned} \sum_{n \geq 1} |\operatorname{sinc}(\pi n s)|^m &\leq \frac{1}{s} \int_0^\infty F_m(x) dx \\ &\leq \frac{1}{s} \left(\sqrt{\frac{3\pi}{8m}} + \frac{\pi^{-m}}{m-1} \right) < \frac{2}{s\sqrt{m}}. \end{aligned}$$

For the last convenient constant, $\pi < 4$ and $\pi^m \geq m$ imply $\sqrt{3\pi/8} + \sqrt{m} \pi^{-m}/(m-1) \leq \sqrt{3/2} + 1/\sqrt{2} < 2$. $\square$

II.313.4 Proof of the uniform lower bound

Fix a positive-gain filter and abbreviate $C = -\log P$, $v = -\log s$, and $R = P/Q > 1$. Suppose first that $R \geq 3/2$. Equation (II.313.14) gives

$$R \leq 1 + \frac{4}{s\sqrt{m}}.$$

Since $R-1 \geq R/3$, it follows that $s \leq 12/(R\sqrt{m})$. Also $4/(s\sqrt{m}) \geq 1/2$, so $1+4/(s\sqrt{m}) \leq 12/(s\sqrt{m})$. Consequently

$$\begin{aligned} v &\geq \log R + \frac{1}{2} \log m - \log 12, \\ C &\geq (m-1)v + \frac{1}{2} \log m - \log 12 \\ &\geq (m-1) \log R + \frac{m}{2} \log m - m \log 12. \end{aligned} \tag{II.313.15}$$

If the last remainder is nonnegative, division by $\log R \leq \log a$ proves the required lower bound, with 20 even replaceable in this case by $2 \log 12$. If it is negative, the baseline $C/\log R \geq m-1$ proves (II.313.5).

Now suppose $1 < R < 3/2$. If $\log m \leq 20$, the correction in (II.313.5) is nonpositive and the baseline suffices. For the remaining case set $x = \log m \geq 20$ and $s_0 = \sqrt{3 \log m/m} < 1$. When $s \leq s_0$,

$$\mathcal{A}_{m,a}(K) \geq \frac{m-1}{2 \log(3/2)} (x - \log x - \log 3). \tag{II.313.16}$$

This is at least $m-1 + mx/(2 \log 2)$. Here is an explicit check of the numerical range. One has

$$\frac{m-1}{m} \geq \frac{99}{100}, \quad \frac{1}{2 \log(3/2)} > \frac{6}{5}, \quad \frac{1}{2 \log 2} < \frac{3}{4}, \quad \log 3 < \frac{11}{10}, \quad \log x \leq \frac{x}{5}.$$

The logarithmic estimates follow respectively from elementary integral or alternating-series bounds and the exponential series; the last one holds at $x = 20$ and persists by differentiation. The left side of (II.313.16), divided by m , is therefore at least $594x/625 - 3267/2500$, whereas the claimed comparison is at most $1 + 3x/4$. The difference is $(501x - 5767)/2500 > 0$ for $x \geq 20$.

If instead $s > s_0$, the right side of (II.313.10) increases with s , and hence

$$\mathcal{A}_{m,a}(K) \geq (m-1) \sqrt{3x/m} m^{2-2/m} \geq \frac{1}{2} \sqrt{3x} m^{3/2}.$$

Since $e^{x/2} \geq x$ for $x \geq 20$, the last expression is at least mx , which is at least $m-1 + mx/(2 \log 2)$. Both subcases dominate (II.313.5) because $a \geq 2$. Together with Lemma II.313.2, this proves both lower estimates in Theorem II.313.1.

II.313.5 Theta filters and their finite approximants

For $t > 0$ put

$$\Theta(t) = \sum_{n \in \mathbb{Z}} e^{-tn^2}, \quad F_t(x) = \sum_{n \in \mathbb{Z}} e^{-tn^2} e(nx).$$

The Gaussian Fourier integral gives the periodization identity

$$F_t(x) = \sqrt{\frac{\pi}{t}} \sum_{j \in \mathbb{Z}} \exp\left(-\frac{\pi^2(x-j)^2}{t}\right) > 0. \quad (\text{II.313.17})$$

For completeness, the n th Fourier coefficient of the right side, after combining the translated intervals into $\mathbb{R}$, is $\sqrt{\pi/t} \int_{\mathbb{R}} e^{-\pi^2 y^2/t} e(-ny) dy = e^{-tn^2}$. Absolute convergence justifies periodization and identifies the Fourier series. At zero this yields

$$\Theta(t) = \sqrt{\frac{\pi}{t}} \Theta(\pi^2/t). \quad (\text{II.313.18})$$

The normalized function $F_t/\Theta(t)$ is nonnegative; its nonnegative Fourier coefficients sum to one, so its values are at most one. It is not a finite polynomial. To remain within (II.313.1), use

$$K_{t,J}(x) = \frac{\sum_{|n| \leq J} \left(1 - \frac{|n|}{J+1}\right) e^{-tn^2} e(nx)}{\sum_{|n| \leq J} \left(1 - \frac{|n|}{J+1}\right) e^{-tn^2}} \quad (J \geq 1). \quad (\text{II.313.19})$$

The numerator is the convolution of F_t with the nonnegative Fejér kernel $(J+1)^{-1} |\sum_{j=0}^J e(jx)|^2$. It is therefore nonnegative. After normalization, its nonnegative even coefficients sum to one, which implies $K_{t,J} \leq 1$ and $K_{t,J}(0) = 1$. Thus every $K_{t,J}$ belongs to the exact finite class $\mathcal{K}$. Its coefficients also decrease with $|n|$, as both factors in the numerator of (II.313.19) do. The upper construction therefore satisfies the stronger decreasing-filter hypotheses of Section II.311 as well.

As $J \rightarrow \infty$, their coefficient sequences converge in ℓ^1 to $e^{-tn^2}/\Theta(t)$ by dominated convergence and convergence of the positive normalizing denominator. The inequality $|u^m - v^m| \leq m|u - v|$ for $u, v \in [0, 1]$ then gives convergence of both P_m and $Q_{m,a}$. The limiting gain is strictly positive, because $a \geq 2$ and the omitted coefficients at $n = \pm 1$ are positive. Hence

$$\mathcal{I}_{m,a} \leq \frac{m \log \Theta(t) - \log \Theta(mt)}{\log \Theta(mt) - \log \Theta(ma^2t)} \quad (t > 0). \quad (\text{II.313.20})$$

This use of a limit proves an upper bound for the infimum over finite polynomials; it does not enlarge the admissible class.

II.313.6 Proof of the uniform upper bound

Put

$$A = \log a, \quad L = \log m, \quad u = L/A, \quad h = \log(e + u), \quad t = \frac{1}{ma^2h}. \quad (\text{II.313.21})$$

In particular $h \geq 1$. For $v \geq 1$, the inequality $n^2 \geq 1 + 3(n-1)$ for $n \geq 1$ gives

$$\Theta(v) \leq 1 + \frac{2e^{-v}}{1 - e^{-3v}} \leq 1 + 3e^{-v}. \quad (\text{II.313.22})$$

Write $\delta = \log \Theta(\pi^2 h)$. Since $\pi^2 > 9$,

$$0 \leq \delta \leq 3e^{-\pi^2 h} \leq 3(e+u)^{-9}. \quad (\text{II.313.23})$$

By (II.313.18), the denominator in (II.313.20) is

$$A + \log \Theta(\pi^2 a^2 h) - \log \Theta(\pi^2 h) \geq A - \delta > 0.$$

The numerator is

$$(m-1)A + \frac{m}{2}L + \frac{m-1}{2}\log(\pi h) + m \log \Theta(\pi^2 m a^2 h) - \log \Theta(\pi^2 a^2 h).$$

Its last positive correction is at most $3me^{-36m} < 1$, by $a^2 \geq 4$, $h \geq 1$, and (II.313.22); its last negative correction can be discarded. Using $\log \pi < 2$ gives the upper bound $(m-1)A + B_\star$, where

$$B_\star = \frac{m}{2}L + \frac{m}{2}\log h + m.$$

Consequently

$$\mathcal{I}_{m,a} - (m-1) \leq \frac{B_\star + (m-1)\delta}{A - \delta}. \quad (\text{II.313.24})$$

We include the elementary constant budget to retain full uniformity in a . Since $A > 1/2$, $e+u > 2$, and $\delta < 3/512 < 1/8$, one has $A - \delta \geq A/2$. Therefore

$$\begin{aligned} \frac{A(B_\star + (m-1)\delta)}{A - \delta} &\leq B_\star + m\delta + \frac{2\delta}{A}(B_\star + m\delta) \\ &\leq B_\star + m[\delta + \delta u + 2\delta \log h + 4\delta(1 + \delta)] \\ &\leq B_\star + m. \end{aligned}$$

For the last line, $\delta \leq 3/512$, $\delta u \leq 3/256$, and $\delta \log h \leq \delta h \leq 3/256$ follow from (II.313.23) and $\log(e+u) \leq e+u$. Also $4\delta(1 + \delta) < 1/32$. The bracket is thus less than $37/512 < 1$. Inserting this bound into (II.313.24) proves

$$\mathcal{I}_{m,a} \leq m - 1 + \frac{(m/2)L + (m/2)\log h + 2m}{A},$$

which is exactly (II.313.6). The proof of Theorem II.313.1 is complete.

II.313.7 Consequences and scope of the cost comparison

Corollary II.313.6 (The arity–torsion regimes). *Let $m \rightarrow \infty$ through integers and let $a = a(m) \geq 2$ be integral.*

- (i) *If $\log a / \log m \rightarrow c \in (0, \infty)$, then $\mathcal{I}_{m,a}/m \rightarrow 1 + 1/(2c)$.*
- (ii) *If $\log a / \log m \rightarrow \infty$, then $\mathcal{I}_{m,a}/m \rightarrow 1$.*
- (iii) *If $\log a = o(\log m)$, then $\mathcal{I}_{m,a} \sim m \log m / (2 \log a)$.*

(iv) For each fixed $a \geq 2$,

$$\mathcal{I}_{m,a} = \frac{m \log m}{2 \log a} + O_a(m(1 + \log \log \log m)).$$

For each fixed $m \geq 2$, one has $\lim_{a \rightarrow \infty} \mathcal{I}_{m,a} = m - 1$.

Proof. Set $u = \log m / \log a$ in Theorem II.313.1. When u stays bounded, the logarithmic correction $\log \log(e + u)$ stays bounded and $\log a \rightarrow \infty$. This proves (i) and (ii). If $u \rightarrow \infty$, then $m - 1 = o(mu)$, and $[1 + \log \log(e + u)] / \log m \rightarrow 0$, because $u \leq \log m / \log 2$. This proves (iii). At fixed a , the upper correction grows as $\log \log \log m$; the lower error is $O_a(m)$, and the term $m - 1$ is absorbed in the stated error. At fixed m , the upper correction tends to zero as $a \rightarrow \infty$, whereas the baseline remains $m - 1$. $\square$

In particular, uniformly whenever $\log a \geq c \log m$ for a fixed $c > 0$, the theorem gives the more precise expression

$$\mathcal{I}_{m,a} = m - 1 + \frac{m \log m}{2 \log a} + O_c\left(\frac{m}{\log a}\right).$$

The limiting coefficient is thus 2 when a is approximately $\sqrt{m}$, $3/2$ when a is approximately m , and $5/4$ when a is approximately m^2 . These coefficients concern the principal filter cost; they are not asserted to be the exponents of a completed Szemerédi bound.

Independent heterogeneous filters do not escape this particular lower bound. If layer ℓ has P_ℓ, Q_ℓ and positive gain, the total principal cost is

$$\frac{\sum_\ell -\log P_\ell}{\sum_\ell \log(P_\ell/Q_\ell)} = \sum_\ell \frac{\log(P_\ell/Q_\ell)}{\sum_j \log(P_j/Q_j)} \mathcal{A}_{m,a}(K_\ell) \geq \mathcal{I}_{m,a}.$$

Zero-gain layers add nonnegative cost and no gain. This is a statement about the specified principal certificate. It gives no lower bound on the survival of a particular input set, which can benefit from extra Fourier relations or structure.

Finally, the finite approximants in (II.313.19) may require a large degree. Their cost estimate alone does not control the torsion term, collision error, or ambient threshold in Theorem II.311.6. The degree-twelve theorem supplies a separate finite certificate where all of those terms are retained. The sharpness of the triple-logarithmic upper correction at fixed a remains open in this manuscript.

II.314 Notation and the main restriction theorem

Source 62, delivered Section 2.

Let $k \geq 2$ and put $m = 2k$. Write

$$\sigma = (\underbrace{1, \dots, 1}_k, \underbrace{-1, \dots, -1}_k).$$

For a subset $B \subseteq \mathbb{Z}/N\mathbb{Z}$, define

$$\mathcal{R}_m(B) = \{(x_1, \dots, x_m) \in B^m : \sigma \cdot x = 0\}, \quad T_m(B) = |\mathcal{R}_m(B)|.$$

Tuples are ordered and repetitions are included. For $\phi : B \rightarrow \mathbb{Z}/M\mathbb{Z}$, let

$$S_m(B, \phi) = |\{x \in \mathcal{R}_m(B) : \sigma \cdot \phi(x) = 0\}|.$$

The nonzero element $\delta(x) = \sigma \cdot \phi(x)$ is the *defect* of an unrespected tuple. We call ϕ a $(1 - \eta)$ -homomorphism of order k on B when

$$S_m(B, \phi) \geq (1 - \eta)T_m(B). \quad (\text{II.314.1})$$

The elementary bound

$$T_m(B) \leq |B|^{m-1} \quad (\text{II.314.2})$$

follows by specifying the first $m - 1$ coordinates and solving for the last.

For $L \geq 2$, define two modulus-independent envelopes:

$$H_{\text{el}}(L) = 1 + 3 \log_2 L, \quad (\text{II.314.3})$$

$$H_{\text{RS}}(L) = 4 + 2 \log(3 \log(2L)). \quad (\text{II.314.4})$$

All unmarked logarithms are natural. The elementary envelope requires no external number-theoretic estimate. The second uses [253, Theorem 15]. Either can be used below, as can their minimum.

Set

$$C_m = \binom{m}{2}, \quad K_m = 1 + \left(\frac{5}{4}\right)^{m-1} 2^{-(m-2)} \left(1 + \frac{2}{m-3}\right) + 2^{4-m}. \quad (\text{II.314.5})$$

In particular,

$$K_{16} < 1.003 < 2, \quad C_{16} = 120. \quad (\text{II.314.6})$$

Theorem II.314.1 (Polynomial-loss restriction for arbitrary cyclic moduli). *Let $N, M \geq 2$, $m = 2k \geq 4$, and $0 < \alpha, \beta, \eta \leq 1$. Suppose $B \subseteq \mathbb{Z}/N\mathbb{Z}$ has cardinality βN and $\phi : B \rightarrow \mathbb{Z}/M\mathbb{Z}$ satisfies*

$$S_m(B, \phi) \geq \alpha |B|^{m-1}. \quad (\text{II.314.7})$$

Choose an integer $L \geq 2$ and an envelope H from (II.314.3)–(II.314.4) such that

$$L \geq \frac{2(m+1)H(L)}{\alpha\eta}. \quad (\text{II.314.8})$$

If

$$N \geq \frac{2(m+1)(K_m + C_m)L^{m-1}}{\alpha\eta\beta^{m-1}}, \quad (\text{II.314.9})$$

then there is $B' \subseteq B$ for which

$$S_m(B', \phi) \geq \frac{\alpha}{(m+1)L^{m-1}} |B|^{m-1}, \quad (\text{II.314.10})$$

$$S_m(B', \phi) \geq (1 - \eta)T_m(B'). \quad (\text{II.314.11})$$

There is no relation required between N and M .

The fact that (II.314.10) counts *respected* tuples is slightly stronger than a lower bound only for all additive tuples. The theorem also implies a direct set-size estimate:

$$|B'| \geq \frac{\alpha^{1/(m-1)}}{(m+1)^{1/(m-1)}L} |B|. \quad (\text{II.314.12})$$

Indeed, combine (II.314.10) with (II.314.2) for B' .

Corollary II.314.2 (Near-prime-scale retention). *For each fixed even $m \geq 4$, there are positive constants c_m, D_m with the following property. Put*

$$x = \alpha\eta, \quad \ell(x) = \log \log(e^e/x).$$

Under (II.314.7), it is sufficient that

$$N \geq D_m \beta^{-(m-1)} x^{-m} \ell(x)^{m-1}$$

to obtain (II.314.11) and

$$S_m(B', \phi) \geq c_m \frac{\alpha^m \eta^{m-1}}{\ell(x)^{m-1}} |B|^{m-1}.$$

Using only H_{el} gives the analogous statement with $\ell(x)$ replaced by $1 + \log(1/x)$.

Proof. For fixed m , an integer L of order $x^{-1}\ell(x)$ satisfies (II.314.8) with H_{RS} . To see this as $x \downarrow 0$, take L to be an integer between $Dx^{-1}\ell(x)$ and twice that quantity. Then $H_{\text{RS}}(L) = (2 + o(1))\ell(x)$, so any sufficiently large constant D depending on m works. Increasing D handles the remaining compact range of $x \in (0, 1]$. Substitute in (II.314.9) and (II.314.10). The elementary version follows in the same way from $H_{\text{el}}(L) = O(1 + \log(1/x))$. $\square$

Theorem II.314.1 is proved in Section II.318. The next three sections establish its analytic and arithmetic ingredients.

II.315 The unit-dilation detector

Source 62, delivered Section 3.

II.315.1 A normalized Fejér filter

On $\mathbb{T} = \mathbb{R}/\mathbb{Z}$, let

$$F_L(t) = \frac{1}{L^2} \left| \sum_{j=0}^{L-1} \exp(2\pi i j t) \right|^2 = \sum_{|c| < L} w_L(c) \exp(2\pi i c t), \quad (\text{II.315.1})$$

$$w_L(c) = \frac{L - |c|}{L^2}, \quad b_L(c) = L w_L(c) = 1 - \frac{|c|}{L} \quad (|c| < L). \quad (\text{II.315.2})$$

Outside this range, both weights are zero. Thus

$$0 \leq F_L \leq 1, \quad w_L(c) \geq 0, \quad \sum_c w_L(c) = 1. \quad (\text{II.315.3})$$

The last equality can also be viewed probabilistically: w_L is the law of $U - V$ for independent uniform $U, V \in \{0, \dots, L-1\}$.

Define the detector kernel and its value at zero by

$$P_{m,L}(t) = \sum_{|c| < L} b_L(c)^m \exp(2\pi i c t), \quad A_{m,L} = P_{m,L}(0). \quad (\text{II.315.4})$$

Although individual cosine terms may be negative, the entire kernel is not.

Lemma II.315.1 (Positivity and central mass). *For every positive integer m , $P_{m,L}(t) \geq 0$. For $m \geq 1$,*

$$A_{m,L} = 1 + 2 \sum_{c=1}^{L-1} (1 - c/L)^m \geq \frac{2L}{m+1}. \quad (\text{II.315.5})$$

Proof. The Fourier coefficients of LF_L are $b_L(c)$. Therefore $P_{m,L}$ is the m -fold convolution of the nonnegative function LF_L with itself, where convolution uses probability Haar measure on $\mathbb{T}$. This proves positivity.

For the second assertion, the function $f(t) = (1 - t/L)^m$ on $[0, L]$ is convex. The trapezoidal rule for a convex function gives

$$\frac{1}{2}f(0) + \sum_{j=1}^{L-1} f(j) + \frac{1}{2}f(L) \geq \int_0^L f(t) dt = \frac{L}{m+1}.$$

Use $f(0) = 1$ and $f(L) = 0$. □

II.315.2 The response of a nonzero defect

Write $U_q = (\mathbb{Z}/q\mathbb{Z})^\times$ and let $\varphi(q) = |U_q|$ be Euler's totient function. Define the primitive response

$$B_{m,L}(q) = \frac{1}{\varphi(q)} \sum_{a \in U_q} P_{m,L}(a/q) \quad (q \geq 2). \quad (\text{II.315.6})$$

For convenience set $B_{m,L}(1) = A_{m,L}$. The Ramanujan sum is

$$c_q(n) = \sum_{a \in U_q} \exp(2\pi i a n / q).$$

The standard orthogonality calculation gives the exact rational formula

$$B_{m,L}(q) = 1 + 2 \sum_{c=1}^{L-1} (1 - c/L)^m \frac{c_q(c)}{\varphi(q)}. \quad (\text{II.315.7})$$

In particular, exact computation of this response does not require numerical trigonometry.

Lemma II.315.2 (Ramanujan identity). *If $g = \gcd(q, n)$, then*

$$\frac{c_q(n)}{\varphi(q)} = \frac{\mu(q/g)}{\varphi(q/g)}. \quad (\text{II.315.8})$$

Also $\varphi(v) \geq \sqrt{v/2}$ for every positive integer v .

Proof. Insert $\mathbf{1}_{\gcd(a,q)=1} = \sum_{d|\gcd(a,q)} \mu(d)$ in the defining sum and apply geometric-sum orthogonality. This yields $c_q(n) = \sum_{d|\gcd(q,n)} d\mu(q/d)$. Evaluating this formula on prime powers, and using multiplicativity, gives (II.315.8).

For the totient bound, write

$$\frac{\varphi(v)^2}{v} = \prod_{p^a \parallel v} p^{a-2}(p-1)^2.$$

Each factor for an odd prime is at least one. The factor at 2, if present, is at least $1/2$. The product is therefore at least $1/2$. □

The relation between this kernel and the restriction procedure is exact. Choose r uniformly in $\mathbb{Z}/N\mathbb{Z}$ and s uniformly in U_M . For each $x \in B$, put

$$p_{r,s}(x) = F_L(rx/N + s\phi(x)/M). \quad (\text{II.315.9})$$

Residue representatives do not affect this expression. For an additive tuple, the Fourier coefficient vectors $c = \lambda\sigma$ contribute $L^{-m}A_{m,L}$ if the tuple is respected, and $L^{-m}B_{m,L}(q)$ if its defect has order q . To justify the latter assertion, write the defect as $(M/q)a$ with $\gcd(a, q) = 1$. The map $U_M \rightarrow U_q$ is surjective with equal-size fibres, so $s\delta/M$ is uniform on the primitive q th roots. Surjectivity follows prime-power by prime-power from the Chinese remainder theorem.

The other coefficient vectors are the arithmetic error. They cannot simply be discarded: averaging over units may give negative Ramanujan sums. Section II.317 bounds their total absolute contribution.

II.316 Uniform primitive-response bounds

Source 62, delivered Section 4.

Lemma II.316.1 (A sampling inequality). *For every $q \geq 2$,*

$$0 \leq B_{m,L}(q) \leq \frac{q}{\varphi(q)}. \quad (\text{II.316.1})$$

Proof. Positivity follows from Lemma II.315.1. Orthogonality on the q -point grid gives

$$\sum_{a=0}^{q-1} P_{m,L}(a/q) = q \left(1 + 2 \sum_{1 \leq j < L/q} (1 - jq/L)^m \right).$$

Since the summand is decreasing in j ,

$$q \sum_{1 \leq j < L/q} (1 - jq/L)^m \leq \int_0^L (1 - t/L)^m dt = \frac{L}{m+1}.$$

Subtracting $P_{m,L}(0) = A_{m,L}$ and using (II.315.5) shows

$$0 \leq \sum_{a=1}^{q-1} P_{m,L}(a/q) \leq q.$$

The primitive points form a subset of these nonzero points, and the terms are nonnegative. Division by $\varphi(q)$ proves the assertion. $\square$

Proposition II.316.2 (Modulus-independent envelopes). *For all $q \geq 2$, $L \geq 2$, and positive integers m ,*

$$B_{m,L}(q) \leq H_{\text{el}}(L). \quad (\text{II.316.2})$$

Using the explicit totient estimate of Rosser–Schoenfeld also gives $B_{m,L}(q) \leq H_{\text{RS}}(L)$.

Proof. First suppose $q \geq L^3$. For $1 \leq c < L$, Lemma II.315.2 gives

$$\left| \frac{c_q(c)}{\varphi(q)} \right| \leq \sqrt{\frac{2 \gcd(q, c)}{q}} \leq \sqrt{\frac{2L}{q}}.$$

Consequently, by (II.315.7),

$$B_{m,L}(q) \leq 1 + 2(L-1)\sqrt{2L/q} < 1 + 2\sqrt{2} < 4. \quad (\text{II.316.3})$$

For $q < L^3$, let $p_1 < \dots < p_t$ be its distinct prime divisors. Since $p_i \geq i+1$,

$$\frac{q}{\varphi(q)} = \prod_{i=1}^t \frac{p_i}{p_i-1} \leq \prod_{i=1}^t \frac{i+1}{i} = t+1 \leq 1 + \log_2 q < 1 + 3 \log_2 L.$$

Combine this with Lemma II.316.1. Because $L \geq 2$, the last expression is at least four, which also covers (II.316.3).

For the sharper envelope, use the classical inequality

$$\frac{q}{\varphi(q)} < e^\gamma \log \log q + \frac{2.50637}{\log \log q} \quad (q \geq 3), \quad (\text{II.316.4})$$

from [253, Theorem 15]. If $q \geq 30$, then $\log \log q > 1$, $e^\gamma < 2$, and the right-hand side is less than $2 \log \log q + 3$. For $q < L^3$ this is at most $2 \log(3 \log L) + 3 < H_{\text{RS}}(L)$. For $2 \leq q < 30$, direct use of the prime divisors gives $q/\varphi(q) \leq 3 < 4$. Finally, (II.316.3) is again below four. $\square$

Remark II.316.3 (The role of the external input). Only the improvement from a single logarithm to a double logarithm uses (II.316.4). The main theorem with H_{el} , the explicit powers in Corollary II.319.2, the exact gcd formula, and the detector obstructions are proved here without it.

II.317 Weighted short relations over composite cyclic groups

Source 62, delivered Section 5.

A union bound over every coefficient vector in $[-L, L]^m$ wastes the normalization $\sum_c w_L(c) = 1$. We retain that normalization throughout. The resulting quantity has an exact arithmetic interpretation.

Let $Z_1, \dots, Z_m$ be independent with distribution w_L , and put

$$p_{L,d}(a) = \mathbb{P}(Z_1 \equiv a \pmod{d}), \quad R_{m,L} = \mathbb{P}(Z_1 = \dots = Z_m) = \sum_c w_L(c)^m.$$

Notice $R_{m,L} = L^{-m} A_{m,L}$; this is a collision probability, not a new filter parameter. For $N \geq 1$, define

$$\kappa_{N,m,L} = \mathbb{E} \left[\gcd(N, Z_1 - Z_m, \dots, Z_{m-1} - Z_m) \mathbf{1}_{\{\exists i < m: Z_i \neq Z_m\}} \right]. \quad (\text{II.317.1})$$

The exclusion of constant coefficient vectors is essential.

Theorem II.317.1 (Exact weighted degeneracy constant). *Let $m \geq 4$ and $L \geq 2$. Then*

$$\kappa_{N,m,L} = \sum_{\substack{d|N \\ 1 \leq d \leq 2L-2}} \varphi(d) \left(\sum_{a \bmod d} p_{L,d}(a)^m - R_{m,L} \right), \quad (\text{II.317.2})$$

$$\kappa_{m,L}^* := \sup_{N \geq 1} \kappa_{N,m,L} = \sum_{d=1}^{2L-2} \varphi(d) \left(\sum_{a \bmod d} p_{L,d}(a)^m - R_{m,L} \right). \quad (\text{II.317.3})$$

The supremum is attained by every multiple of $\text{lcm}(1, 2, \dots, 2L-2)$. Moreover,

$$0 \leq \kappa_{N,m,L} \leq \kappa_{m,L}^* \leq K_m, \quad (\text{II.317.4})$$

where K_m is given in (II.314.5), and

$$\lim_{L \rightarrow \infty} \kappa_{m,L}^* = \sum_{d=1}^{\infty} \frac{\varphi(d)}{d^{m-1}} = \frac{\zeta(m-2)}{\zeta(m-1)}. \quad (\text{II.317.5})$$

Proof. For a nonconstant tuple, the gcd of its integer differences is a positive integer at most $2L-2$. Apply $g = \sum_{d|g} \varphi(d)$ to the gcd in (II.317.1). The event that all differences are divisible by d is the event that all Z_i have the same residue modulo d , with probability $\sum_a p_{L,d}(a)^m$. Subtract the probability $R_{m,L}$ of an actually constant tuple. This proves (II.317.2). All bracketed terms are nonnegative. Dropping $d | N$ gives (II.317.3), and making N a multiple of every $d \leq 2L-2$ achieves equality.

We next bound the residue probabilities. Write $L = ud + r$ with $0 \leq r < d$. The residue counts of a uniform variable in $\{0, \dots, L-1\}$ are u or $u+1$. Since Z_1 is a difference of two such variables, Cauchy–Schwarz shows that its largest residue probability is attained at zero. Hence

$$\max_a p_{L,d}(a) = p_{L,d}(0) = \frac{1}{d} + \frac{r(d-r)}{dL^2} \leq \frac{1}{d} + \frac{d}{4L^2}. \quad (\text{II.317.6})$$

For $d \leq L$ this is at most $5/(4d)$. For $d > L$, the original uniform variable has L distinct residues, and $p_{L,d}(0) = 1/L$.

The term $d = 1$ in (II.317.3) is at most one. Use $\sum_a p_a^m \leq (\max_a p_a)^{m-1}$ for the other terms, and discard the nonnegative subtraction $R_{m,L}$ to get

$$\begin{aligned} \kappa_{m,L}^* &\leq 1 + \left(\frac{5}{4}\right)^{m-1} \sum_{d=2}^L \frac{\varphi(d)}{d^{m-1}} + L^{-(m-1)} \sum_{d=L+1}^{2L-2} \varphi(d) \\ &\leq 1 + \left(\frac{5}{4}\right)^{m-1} \sum_{d=2}^{\infty} d^{-(m-2)} + 2L^{3-m}. \end{aligned}$$

The integral estimate

$$\sum_{d=2}^{\infty} d^{-(m-2)} \leq 2^{-(m-2)} + \int_2^{\infty} t^{-(m-2)} dt = 2^{-(m-2)} \left(1 + \frac{2}{m-3}\right)$$

and $L \geq 2$ prove (II.317.4).

For the limit, fix d . Formula (II.317.6), or the underlying residue counts, shows $p_{L,d}(a) \rightarrow 1/d$. Also $R_{m,L} = O_m(L^{1-m}) \rightarrow 0$. For $d \leq L$, the summands are bounded by a constant depending on m times $d^{-(m-2)}$, a summable sequence. The contribution from $d > L$ is at most $2L^{3-m} \rightarrow 0$. Truncation at a fixed d , followed by the summable tail bound, proves the first equality of (II.317.5). The last equality follows by multiplying the absolutely convergent Euler products, or from $\varphi(n) = n \sum_{d|n} \mu(d)/d$. $\square$

II.317.1 Why the gcd is the correct finite-ring count

Lemma II.317.2 (Two linear equations). *For $c \in \mathbb{Z}^m$ not an integer multiple of σ , the number of $x \in (\mathbb{Z}/N\mathbb{Z})^m$ satisfying $\sigma \cdot x = c \cdot x = 0$ equals*

$$N^{m-2} \gcd(N, d_1, \dots, d_{m-1}), \quad d_i = c_i - c_m \sigma_i / \sigma_m. \quad (\text{II.317.7})$$

Consequently,

$$\sum_{c \notin \mathbb{Z}\sigma} \left(\prod_{j=1}^m w_L(c_j) \right) \# \{x : \sigma \cdot x = c \cdot x = 0\} = N^{m-2} \kappa_{N,m,L}. \quad (\text{II.317.8})$$

Proof. Because $\sigma_m = \pm 1$, solve the first equation for x_m . The second becomes $\sum_{i < m} d_i x_i = 0$. The image of this homomorphism from $(\mathbb{Z}/N\mathbb{Z})^{m-1}$ to $\mathbb{Z}/N\mathbb{Z}$ has size N/g , where $g = \gcd(N, d_1, \dots, d_{m-1})$. Its kernel therefore has size $N^{m-1}/(N/g) = gN^{m-2}$. This uses no field structure.

Multiply each coefficient c_i by σ_i . The distribution is unchanged because w_L is symmetric. The differences in the gcd become differences from the last transformed coefficient, up to sign, and $c \in \mathbb{Z}\sigma$ becomes the event that all transformed coefficients are equal. Averaging (II.317.7) gives (II.317.8). $\square$

II.317.2 Exact small lengths and a false limiting-constant shortcut

Proposition II.317.3. *For every integer $m \geq 4$,*

$$\kappa_{m,2}^* = 1 - 4^{1-m}, \quad (\text{II.317.9})$$

$$\kappa_{m,3}^* = 1 + \frac{5^m + 4^m + 2 \cdot 3^m - 6 \cdot 2^m - 12}{9^m}. \quad (\text{II.317.10})$$

In particular, at $m = 16$ the limit in (II.317.5) is not a valid bound for every finite L .

Proof. For $L = 2$, the weights at $-1, 0, 1$ are $1/4, 1/2, 1/4$. A nonconstant tuple has integer gcd two exactly when all its entries are in $\{-1, 1\}$ and both signs occur. All other nonconstant tuples have gcd one. This gives (II.317.9).

For $L = 3$, the integer weight numerators are $(1, 2, 3, 2, 1)$, with denominator nine. The residue-power numerators for $d = 1, 2, 3, 4$ are respectively

$$9^m, \quad 5^m + 4^m, \quad 3 \cdot 3^m, \quad 3^m + 3 \cdot 2^m.$$

The constant-tuple numerator is $3^m + 2 \cdot 2^m + 2$. Insert these into (II.317.3) with totient weights $1, 1, 2, 2$.

Finally,

$$\frac{\zeta(14)}{\zeta(15)} = 1 + \sum_{d=2}^{\infty} \frac{\varphi(d)}{d^{15}} \leq 1 + 2^{-14} \left(1 + \frac{2}{13} \right) < \kappa_{16,3}^*.$$

The last inequality follows directly from the integer expression (II.317.10); it is also checked as a rational inequality by the accompanying verifier. $\square$

This distinction matters when inserting constants into a formal argument. The exact asymptotic coefficient is not automatically a uniform finite-size coefficient. In decimal form, $\kappa_{16,3}^* \approx 1.0000847096$, whereas $\zeta(14)/\zeta(15) \approx 1.0000306590$. The safe rational envelope (II.314.5) gives $K_{16} \approx 1.0022457446$.

II.318 The restriction proof, including repeated vertices

Source 62, delivered Section 6.

For a family $\mathcal{Q} \subseteq \mathcal{R}_m(B)$, define its weighted count at fixed r, s by

$$W_{\mathcal{Q}}(r, s) = \sum_{x \in \mathcal{Q}} \prod_{j=1}^m p_{r,s}(x_j).$$

The product here has one factor per *position*, even when coordinates repeat. This convention is useful for Fourier expansion, but differs from actual subset selection. We keep the two operations separate.

Lemma II.318.1 (Weighted main term and error). *Let G be the respected tuple family in B , and let Y_q count the unrespected tuples whose defect has order q . Then*

$$|\mathbb{E}W_G - L^{-m}A_{m,L}S_m(B, \phi)| \leq \kappa_{N,m,L}N^{m-2}, \quad (\text{II.318.1})$$

$$\left| \mathbb{E}W_{\mathcal{R}_m(B) \setminus G} - L^{-m} \sum_{q|M, q>1} B_{m,L}(q)Y_q \right| \leq \kappa_{N,m,L}N^{m-2}. \quad (\text{II.318.2})$$

Proof. Expand the product using (II.315.1) and average over r . A coefficient vector c survives this average exactly when $c \cdot x = 0$ in $\mathbb{Z}/N\mathbb{Z}$. Vectors $c = \lambda\sigma$ yield the main terms calculated in Section II.315.

For every other c , the remaining average over s is a complex number of absolute value at most one. Thus its absolute contribution, summed over any subfamily of additive tuples, is bounded by

$$\sum_{c \notin \mathbb{Z}\sigma} \left(\prod_j w_L(c_j) \right) \#\{x \in (\mathbb{Z}/N\mathbb{Z})^m : \sigma \cdot x = c \cdot x = 0\}.$$

Apply Lemma II.317.2. The argument is valid separately for the good and the bad subfamilies. In particular, no positivity of the individual Ramanujan coefficients has been assumed. $\square$

Now, conditional on r, s , include each *distinct point* $x \in B$ in B' independently with probability $p_{r,s}(x)$. Let G' and Y' be the actual respected and unrespected tuple counts in B' . A repeated point is included once, not independently once per position.

Lemma II.318.2 (Rounding correction). *The above procedure satisfies*

$$\mathbb{E}G' \geq L^{-m}A_{m,L}S_m(B, \phi) - \kappa_{N,m,L}N^{m-2}, \quad (\text{II.318.3})$$

$$\mathbb{E}Y' \leq L^{-m} \sum_{q|M, q>1} B_{m,L}(q)Y_q + (\kappa_{N,m,L} + C_m)N^{m-2}. \quad (\text{II.318.4})$$

Proof. For a tuple with distinct support D , its conditional survival probability is $\prod_{x \in D} p_{r,s}(x)$. Because $0 \leq p_{r,s}(x) \leq 1$, this is at least the position-by-position product. This proves the favourable inequality for the good family.

For a tuple with no repetitions the two products agree. Their difference is at most one otherwise. For any fixed pair of distinct positions i, j , the number of additive tuples satisfying $x_i = x_j$ is N^{m-2} : identify the two variables and solve the additive relation for one of the other variables, whose coefficient is still 1 or -1 . Such another variable exists because $m \geq 4$. A union bound gives at most $C_m N^{m-2}$ repeated tuples. Combine this with Lemma II.318.1. $\square$

The following theorem is the full data-sensitive form of the method.

Theorem II.318.3 (Restriction retaining the defect-order profile). *Fix $L \geq 2$ and $0 < \eta \leq 1$. Write $S = S_m(B, \phi)$, and let Y_q be as above. Define*

$$\Pi = L^{-m} \left(\eta A_{m,L} S - (1 - \eta) \sum_{q|M, q>1} B_{m,L}(q) Y_q \right), \quad (\text{II.318.5})$$

$$\Delta = (\kappa_{N,m,L} + (1 - \eta) C_m) N^{m-2}. \quad (\text{II.318.6})$$

If $\Pi > \Delta$, there is $B' \subseteq B$ satisfying (II.314.11) and

$$S_m(B', \phi) \geq \frac{\Pi - \Delta}{\eta}.$$

Proof. Use the linear objective

$$Z = \eta G' - (1 - \eta) Y' = \eta(G' + Y') - Y'. \quad (\text{II.318.7})$$

Lemmas II.318.1–II.318.2 imply $\mathbb{E}Z \geq \Pi - \Delta > 0$. Some realization therefore has $Z \geq \Pi - \Delta$. The inequality $Z \geq 0$ is exactly $Y' \leq \eta(G' + Y')$. Also $Z \leq \eta G'$, so $G' \geq (\Pi - \Delta)/\eta$. Both assertions follow. $\square$

Proof of Theorem II.314.1. Put $b = |B|$ and $S = S_m(B, \phi)$. By (II.314.2), $\sum_q Y_q \leq b^{m-1}$. Proposition II.316.2 and Lemma II.318.2, applied to (II.318.7), give

$$\mathbb{E}Z \geq L^{-m} (\eta \alpha A_{m,L} - (1 - \eta) H(L)) b^{m-1} - (K_m + C_m) N^{m-2}.$$

By (II.315.5) and (II.314.8), $\alpha \eta A_{m,L} \geq 4H(L)$. Hence the first term is at least $\frac{3}{4} \eta \alpha A_{m,L} L^{-m} b^{m-1}$. On the other hand, (II.314.9), $b = \beta N$, and (II.315.5) give

$$(K_m + C_m) N^{m-2} \leq \frac{\alpha \eta}{2(m+1)L^{m-1}} b^{m-1} \leq \frac{1}{4} \eta \alpha A_{m,L} L^{-m} b^{m-1}.$$

It follows that

$$\mathbb{E}Z \geq \frac{1}{2} \eta \alpha A_{m,L} L^{-m} b^{m-1} > 0.$$

Choose a realization with at least this value. As in Theorem II.318.3, it has the required quality and

$$G' \geq \frac{1}{2} \alpha A_{m,L} L^{-m} b^{m-1} \geq \frac{\alpha}{(m+1)L^{m-1}} b^{m-1}.$$

$\square$

Remark II.318.4 (No exceptional-set deletion). The proof never asserts that every selected tuple is nondegenerate. It controls the total signed error from all extra short relations and then separately controls repeated vertices. This is why the filter length does not appear in the constant K_m .

II.319 An explicit order-eight strengthening

Source 62, delivered Section 7.

The main theorem can be used with exact integer parameter choices. For $x = \alpha\eta$, let $e \geq 1$ be the least integer satisfying

$$2^e \geq \frac{34(1+3e)}{x}, \quad L = 2^e. \quad (\text{II.319.1})$$

Such an integer exists because exponential growth dominates linear growth.

Corollary II.319.1 (A rationally checkable parameter recipe). *Under the hypotheses of Theorem II.314.1 with $m = 16$, the sufficient condition*

$$N \geq \frac{4148L^{15}}{\alpha\eta\beta^{15}} \quad (\text{II.319.2})$$

with L from (II.319.1) gives a subset B' satisfying

$$S_{16}(B', \phi) \geq \frac{\alpha}{17L^{15}}|B|^{15}, \quad S_{16}(B', \phi) \geq (1-\eta)T_{16}(B').$$

Proof. At dyadic L , $H_{\text{el}}(L) = 1 + 3e$, so (II.319.1) is exactly (II.314.8). Use $K_{16} < 2$, $C_{16} = 120$, and $2 \cdot 17 \cdot 122 = 4148$ in (II.314.9). $\square$

Corollary II.319.2 (Fully explicit polynomial powers). *Let $0 < \alpha, \beta, \eta \leq 1$, let $B \subseteq \mathbb{Z}/N\mathbb{Z}$ have size βN , and let $\phi : B \rightarrow \mathbb{Z}/M\mathbb{Z}$ respect at least $\alpha|B|^{15}$ additive sixteen-tuples. If*

$$N \geq 2^{253}(\alpha\eta)^{-31}\beta^{-15}, \quad (\text{II.319.3})$$

then some $B' \subseteq B$ satisfies

$$S_{16}(B', \phi) \geq 2^{-245}\alpha^{31}\eta^{30}|B|^{15}, \quad S_{16}(B', \phi) \geq (1-\eta)T_{16}(B'). \quad (\text{II.319.4})$$

Proof. Again write $x = \alpha\eta$ and put $y = 136/x$. Choose an integer $j = \lceil 2 \log_2 y \rceil$, so that $y^2 \leq 2^j < 2y^2$ and $j \leq 2 \log_2 y + 1$. For $y \geq 4$,

$$1 + \frac{3}{2} \log_2 y \leq y;$$

this holds at $y = 4$ and the difference of the right and left sides is increasing thereafter. Thus

$$\frac{34(1+3j)}{x} \leq y(1 + \frac{3}{2} \log_2 y) \leq y^2 \leq 2^j.$$

Minimality in (II.319.1) gives

$$L \leq 2^j < 2(136/x)^2 < 2^{16}x^{-2}. \quad (\text{II.319.5})$$

Consequently,

$$\frac{\alpha}{17L^{15}} > 2^{-245}\alpha x^{30} = 2^{-245}\alpha^{31}\eta^{30},$$

because $17 < 2^5$. Also

$$\frac{4148L^{15}}{x\beta^{15}} < 2^{253}x^{-31}\beta^{-15},$$

because $4148 < 2^{13}$. Apply Corollary II.319.1. $\square$

[write] **7 October 2026; sources 60–69.** *Checked at the write.* With $M = N$, $|B| = \beta N$ and $0 < \alpha, \eta \leq 1$: $(\alpha\eta/4)^{2^{19}} \leq 2^{-2^{20}}(\alpha\eta)^{2^{19}} \leq 2^{-245}\alpha^{31}\eta^{30}$, so this corollary implies lemma_9_3 with $N_0 = \lceil 2^{253}(\alpha\eta)^{-31}\beta^{-15} \rceil$, on paper. Against source 68's exponent-53 theorem, whose retention of respected tuples is at least $(1-\eta)(\alpha\eta/4)^{53}\beta^{15}$, this corollary retains more exactly when $\alpha^{22}\eta^{23}(1-\eta) < 2^{-139}$; its threshold grows like β^{-15} , source 68's like β^{-1} (Section II.310).

II.319.1 Comparison with the repository interface

The formal interface `lemma_9_3` asks for a retained additive count

$$(\alpha\eta/4)^{2^{19}} \beta^{15} N^{15} \tag{II.319.6}$$

and a $(1 - \eta)$ -homomorphism of order eight [255]. Our new lower bound is larger throughout the nonvacuous parameter range. Indeed, since $x = \alpha\eta \leq 1$ and $\alpha \geq x$,

$$2^{-245} \alpha^{31} \eta^{30} = 2^{-245} \alpha x^{30} \geq 2^{-245} x^{31} \geq (x/4)^{2^{19}}.$$

The assertion for $\alpha > 1$ is vacuous by (II.314.2); similarly, $\beta > 1$ cannot describe a subset of $\mathbb{Z}/N\mathbb{Z}$. Thus the written theorem implies the full old interface, with an explicit size threshold, and considerably stronger retention.

The coarse powers 31, 30 are not the strongest output. They are useful because they contain no implicit scale choice and use only the elementary response envelope. The data-sensitive theorem and Corollary II.314.2 retain the near-prime exponents 16, 15, at the cost of a double-logarithmic denominator.

II.319.2 Numerical scale of the sufficient conditions

Table II.319.1 uses only the elementary, dyadic recipe, with $\beta = 1/2$. The entries are logarithms of the retained coefficient $\alpha/(17L^{15})$ and the sufficient threshold in (II.319.2). They are evaluations of proved formulas, not optimized finite examples. Even the improved sufficient thresholds remain very large.

α	η	e	L	$\log_{10}$ retention	$\log_{10} N_0$	$\log_{10}$ old retention
1/2	1/2	13	8192	−60.23	67.44	−631305.66
1/10	1/10	18	262144	−83.51	91.41	−1364228.83
1/100	1/100	25	33554432	−116.12	125.02	−2412804.83

Table II.319.1: All-modulus sufficient estimates for additive sixteen-tuples. The old coefficient is (II.319.6) without $\beta^{15} N^{15}$.

II.320 Torsion obstructions and the scope of optimality

Source 62, delivered Section 8.

II.320.1 Why averaging all scalars is insufficient

Suppose the target dilation s were uniform on all of $\mathbb{Z}/M\mathbb{Z}$, rather than its unit group. For a defect of order q , the proportional-vector response would be

$$B_{m,L}^{\text{all}}(q) = \sum_{\substack{|c| < L \\ q|c}} b_L(c)^m. \tag{II.320.1}$$

For fixed q and m , Riemann sums show

$$A_{m,L} = \frac{2L}{m+1} + O_m(1), \quad B_{m,L}^{\text{all}}(q) = \frac{2L}{q(m+1)} + O_m(1).$$

Therefore

$$\lim_{L \rightarrow \infty} \frac{A_{m,L}}{B_{m,L}^{\text{all}}(q)} = q. \quad (\text{II.320.2})$$

For an order-two defect the gain saturates at two. A filter-length increase alone cannot provide the arbitrarily large gain needed as $\alpha\eta \rightarrow 0$. This is a failure of that averaging calculation, not a counterexample to the restriction theorem.

In contrast, the unit response for each fixed $q > 1$ tends to zero. Here is a useful quantitative estimate that makes this precise.

Lemma II.320.1 (Suppression of a fixed torsion order). *For $t \notin \mathbb{Z}$,*

$$P_{m,L}(t) \leq \frac{m}{L \sin^2(\pi t)}.$$

Consequently, for fixed m, q , $B_{m,L}(q) = O_{m,q}(L^{-1})$.

Proof. Let $a_n = (1 - n/L)^m$ for $0 \leq n \leq L$, and $a_n = 0$ afterwards. Convexity gives $d_n = a_{n-1} - 2a_n + a_{n+1} \geq 0$ for $n \geq 1$, with $\sum_{n \geq 1} d_n = a_0 - a_1$. Multiplication of the symmetric Fourier series by $2 - 2\cos(2\pi t)$ gives

$$4\sin^2(\pi t)P_{m,L}(t) = 2(a_0 - a_1) - 2 \sum_{n \geq 1} d_n \cos(2\pi nt) \leq 4(a_0 - a_1) \leq 4m/L.$$

Average over the finite set of primitive q th roots. □

II.320.2 A necessary double logarithm for this detector

The fixed- q estimate cannot be made uniform with a constant independent of q and L .

Theorem II.320.2 (Primorial lower bound for the primitive response). *For every fixed positive integer m , there is an unbounded sequence of integers L such that*

$$\sup_{q \geq 2} B_{m,L}(q) \geq \frac{\log \log(6L)}{6(m+1)}. \quad (\text{II.320.3})$$

Together with Proposition II.316.2, this identifies the worst-case double-logarithmic order, up to constants, along that sequence.

Proof. Let Q be the product of all primes at most an integer $y \geq 5$, and put $L = Q/6$. The residues 1 and $Q - 1$ are units. For $|c| < L = Q/6$, $\cos(2\pi c/Q) \geq 1/2$, so

$$P_{m,L}(1/Q) \geq \frac{1}{2}A_{m,L}.$$

By positivity, the two primitive residues just mentioned give

$$B_{m,L}(Q) \geq \frac{A_{m,L}}{\varphi(Q)} \geq \frac{Q}{3(m+1)\varphi(Q)}.$$

No prime number theorem is needed for the remaining estimate. The Euler product

$$\frac{Q}{\varphi(Q)} = \prod_{p \leq y} (1 - p^{-1})^{-1}$$

expands as the reciprocal sum over positive integers all of whose prime factors are at most y . It includes every integer up to y , whence $Q/\varphi(Q) \geq \sum_{n=1}^y 1/n \geq \log y$. Also $Q \leq y! \leq y^y$, and therefore $\log \log Q \leq \log y + \log \log y \leq 2 \log y$. Combining the last three inequalities yields (II.320.3). Infinitely many primes ensure that these L are unbounded. $\square$

Remark II.320.3 (What has, and has not, been shown sharp). Theorem II.320.2 rules out an absolute bound for $B_{m,L}(q)$ for the specific kernels (II.315.4) and primitive-unit averaging. It does not prove that every restriction theorem needs a double logarithm. A different kernel, several coordinated tests, or a structural restriction argument might avoid it. Nor does the theorem assert sharpness of the powers of α and η in Theorem II.314.1.

II.321 Extensions and constructive content

Source 62, delivered Section 9.

II.321.1 Different domain and target moduli

The whole proof allows $\mathbb{Z}/N\mathbb{Z}$ as domain and $\mathbb{Z}/M\mathbb{Z}$ as target. The character detecting a domain relation is chosen independently from the unit dilation detecting a target defect. Domain degeneracy depends only on N , while the defect response depends only on the order of the target error. Thus neither coprimality of N, M nor a bound on the target order is needed. A trivial target has no bad tuples and can be handled by $B' = B$.

This separation is useful when an auxiliary cyclic group records a discrete phase or label but the original configuration lives in another cyclic group. It is not an assertion about general noncyclic targets: those do not possess a single unit orbit with the same primitive-character property.

II.321.2 Labelled domains and multifunctions

Let X be a finite set of labels, let $r : X \rightarrow \mathbb{Z}/N\mathbb{Z}$ be a base map, and let $\phi : X \rightarrow \mathbb{Z}/M\mathbb{Z}$ be a target map. Suppose $|r^{-1}(a)| \leq R$ for every a , where $R \geq 1$ is an integer, and $|X| = \beta RN$. Define additive and respected tuples using $r(x_j)$ in the domain equation and $\phi(x_j)$ in the target equation. These are tuples of labels, not unlabelled base points.

Theorem II.321.1 (Bounded-fibre labelled restriction). *Let $m = 2k \geq 4$, $N, M \geq 2$, and $0 < \alpha, \beta, \eta \leq 1$. Suppose at least $\alpha R |X|^{m-1}$ labelled additive tuples are respected. Choose L as in (II.314.8). It is sufficient that*

$$N \geq \frac{2(m+1)(K_m + C_m/R)L^{m-1}}{\alpha\eta\beta^{m-1}} \quad (\text{II.321.1})$$

to obtain $X' \subseteq X$ that respects at least a proportion $1 - \eta$ of its labelled additive tuples and contains at least

$$\frac{\alpha}{(m+1)L^{m-1}} R |X|^{m-1}$$

respected labelled tuples.

Proof. The additive tuple cap is now $R|X|^{m-1}$, since the last base coordinate is determined but has at most R labels. Apply the filter

$$F_L(ur(x)/N + s\phi(x)/M)$$

and round independently for each label.

For a fixed base tuple satisfying two linear equations, there are at most R^m label lifts. Thus the error in Lemma II.318.1 is at most $\kappa_{N,m,L}R^mN^{m-2}$. For a specified equality of two labels, the number of additive tuples is at most $R^{m-1}N^{m-2}$: after identifying the labels, solve the base equation for one remaining label, and use the fibre cap. Hence the rounding correction is at most $C_mR^{m-1}N^{m-2}$.

The objective (II.318.7) therefore has error at most $(K_m + C_m/R)R^mN^{m-2}$. The main scale is $R|X|^{m-1} = \beta^{m-1}R^mN^{m-1}$. The proof of Theorem II.314.1, with these substitutions, gives (II.321.1) and the claimed conclusion. $\square$

The normalization by R is necessary. Omitting it would incorrectly use the graph bound (II.314.2) in a setting where several labels can have the same base coordinate. This version fits the multifunction viewpoint used later in Gowers' argument, but by itself supplies no statement about higher-dimensional arrangements.

II.321.3 Finite conditional-expectation construction

The existence proof is constructive in a finite, though not necessarily practical, sense. First enumerate the pairs $(r, s) \in \mathbb{Z}/N\mathbb{Z} \times U_M$ and choose one whose conditional expectation of (II.318.7) is at least the unconditional mean. Then decide membership in B' one point at a time. Conditional on all previous decisions, the remaining objective has an expectation affine in the next membership indicator, so at least one of the decisions zero or one does not decrease that expectation.

All filter values belong to a cyclotomic number field and are real algebraic numbers, so the required comparisons can in principle be made exactly. The proof does not establish an attractive bit-complexity bound for this implementation. In particular, it should not be confused with a fast algorithm for finding long arithmetic progressions.

II.322 Notation and the exact relation interface

Source 61, delivered Section 2.

For subsets of an abelian group, $A + B = \{a + b : a \in A, b \in B\}$ and $A - B$ is defined similarly. The set mA is the m -fold sumset; repetitions are permitted, and $0A = \{0\}$. Unless otherwise specified, all sums over group elements use counting measure. For nonempty A , translating by a fixed element shows $|A - A| \geq |A|$, so every difference ratio below is at least one.

Definition II.322.1 (Freiman homomorphism of order m). A map $\phi : B \rightarrow H$ is a Freiman homomorphism of order m if, for all $x_1, \dots, x_m, y_1, \dots, y_m \in B$,

$$\sum_{i=1}^m x_i = \sum_{i=1}^m y_i \implies \sum_{i=1}^m \phi(x_i) = \sum_{i=1}^m \phi(y_i).$$

Only preservation is required; injectivity of ϕ is not required.

Repetitions are essential. Removing repeated variables changes the statement and destroys the torsion obstructions below. If B is nonempty, an order- m homomorphism is an order- s homomorphism for every $1 \leq s \leq m$: append $m - s$ copies of a fixed element to both sides.

Definition II.322.2 (Relation energies). For $s \geq 1$ and $A \subseteq G$, set

$$\mathcal{E}_s(A) = \#\{(x_1, \dots, x_s, y_1, \dots, y_s) \in A^{2s} : \sum_i x_i = \sum_i y_i\}.$$

For $\phi : B \rightarrow H$, set

$$\mathcal{E}_s^\phi(B) = \#\{(\mathbf{x}, \mathbf{y}) \in B^{2s} : \sum_i x_i = \sum_i y_i, \sum_i \phi(x_i) = \sum_i \phi(y_i)\}.$$

Equivalently, $\mathcal{E}_s^\phi(B) = \mathcal{E}_s(\Gamma)$.

The tuples are ordered and may be degenerate. Specifying the first $2s - 1$ entries determines the last, if it belongs to the set. Hence

$$\mathcal{E}_s^\phi(B) \leq \mathcal{E}_s(B) \leq n^{2s-1}. \quad (\text{II.322.1})$$

If ϕ is exact of order s , the first inequality is equality.

For functions $u, v : G \rightarrow \mathbb{C}$, use convolution

$$(u * v)(z) = \sum_{x \in G} u(x)v(z - x).$$

If w is supported on A , define

$$\mathcal{E}_s(w) = \sum_{z \in G} |w^{*s}(z)|^2. \quad (\text{II.322.2})$$

This equals the corresponding sum over equal-sum tuples, weighted by $\prod_i w(x_i) \overline{w(y_i)}$. For $w = \mathbf{1}_A$ it is $\mathcal{E}_s(A)$. The graph weight $\tilde{w}$ is defined by $\tilde{w}(x, \phi(x)) = w(x)$ and zero elsewhere.

Lemma II.322.3 (Weighted exactness). *If $\phi|_A$ is a Freiman homomorphism of order m , then*

$$\mathcal{E}_s(w) = \mathcal{E}_s(\tilde{w}) \quad (1 \leq s \leq m)$$

for every complex weight w supported on A .

Proof. If A is empty both energies vanish. Otherwise expand both energies into ordered pairs of s -tuples. The graph sum is subject to the additional equality of image sums. Exactness and padding show that this additional condition holds for every base relation. The indexing sets and their weights therefore agree term by term. $\square$

II.323 Separating every nonzero defect with few characters

Source 61, delivered Section 3.

Write $\mathbb{T} = \mathbb{R}/\mathbb{Z}$, and let $\|t\|_{\mathbb{T}}$ be the distance to zero. A character is an additive homomorphism $\chi : H \rightarrow \mathbb{T}$. Its exponential form is $\exp(2\pi i \chi(x))$. The finite dual group is denoted $\hat{H}$.

Lemma II.323.1 (A uniform root-grid estimate). *Let $h \neq 0$ have order q in a finite abelian group H . For a uniformly chosen $\chi \in \widehat{H}$,*

$$\Pr(\|\chi(h)\|_{\mathbb{T}} < 1/5) = \frac{2 \lceil q/5 \rceil - 1}{q} \leq \frac{1}{2}. \quad (\text{II.323.1})$$

Proof. Evaluation at h maps $\widehat{H}$ onto the cyclic group of q -th roots in $\mathbb{T}$. One way to check surjectivity is to write H as a product of cyclic groups. The possible evaluations in each coordinate form a cyclic subgroup of order equal to the order of that coordinate of h ; together they generate a subgroup of order q . Every fiber of a surjective homomorphism between finite groups has the same size. Thus the evaluation is uniform on $\{j/q : 0 \leq j < q\}$.

A residue has distance strictly less than $1/5$ precisely when its signed representative has absolute value strictly less than $q/5$. There are $2 \lceil q/5 \rceil - 1$ such residues. For $q \geq 10$,

$$2 \lceil q/5 \rceil - 1 < 2q/5 + 1 \leq q/2.$$

For $q = 2, 3, 4, 5, 6, 7, 8, 9$ the numerators are, respectively, 1, 1, 1, 1, 3, 3, 3, which also satisfy the required inequality. $\square$

[write] **7 October 2026; sources 60–69.** *A misprint in the proof* (standing rule of 4 October 2026; checked at the write). The list “1, 1, 1, 1, 3, 3, 3” has seven numerators for the eight orders $q = 2, \dots, 9$; the values of $2 \lceil q/5 \rceil - 1$ are 1, 1, 1, 1, 3, 3, 3, the last ($q = 9$) missing. All eight satisfy $2 \lceil q/5 \rceil - 1 \leq q/2$ (with equality at $q = 2$ and $q = 6$), so the lemma holds as stated.

The strict inequality at radius $1/5$ is deliberate. Replacing it by a closed ball would include three of the five roots when $q = 5$, invalidating the one-half bound. No large-order approximation is being used.

Lemma II.323.2 (Greedy separation). *For $D \subseteq H \setminus \{0\}$ there are characters $\chi_1, \dots, \chi_r$ with*

$$r \leq \lceil \log_2(|D| + 1) \rceil$$

such that every $d \in D$ satisfies $\|\chi_j(d)\|_{\mathbb{T}} \geq 1/5$ for at least one j .

Proof. If D is empty, take no characters. Otherwise, apply Lemma II.323.1 and linearity of expectation to the current unseparated set. Some character leaves at most half that set inside the open radius- $1/5$ ball. Choose such a character and continue. After r choices the remaining number is at most $|D|2^{-r}$. For $r = \lceil \log_2(|D| + 1) \rceil$ this is strictly less than one, hence zero. The argument is a finite greedy construction; independence between the defects is unnecessary. $\square$

Theorem II.323.3 (Universal higher-order avoidance coloring). *Let H be a finite abelian group, $D \subseteq H \setminus \{0\}$, and $m \geq 1$. There is a partition of H into at most*

$$(5m)^{\lceil \log_2(|D|+1) \rceil} \leq 5m(|D| + 1)^{\log_2(5m)} \quad (\text{II.323.2})$$

nonempty cells S such that $(mS - mS) \cap D = \emptyset$. Symmetry of D is not needed.

Proof. Choose the characters in Lemma II.323.2, and represent every character value in $[0, 1)$. Partition $[0, 1)$ into the $5m$ half-open intervals

$$I_a = [a/(5m), (a+1)/(5m)), \quad 0 \leq a < 5m.$$

Color $x \in H$ by the tuple of intervals containing its character values. There are at most $(5m)^r$ colors.

Suppose $x_1, \dots, x_m, y_1, \dots, y_m$ lie in one color cell. For every chosen character, all $2m$ representatives lie in an interval of length $1/(5m)$, so the difference of the two real sums has absolute value strictly less than $1/5$. Consequently

$$\left\| \chi_j \left(\sum_i x_i - \sum_i y_i \right) \right\|_{\mathbb{T}} < 1/5 \quad \text{for every } j.$$

This difference cannot belong to D , by separation. The final inequality in (II.323.2) follows from $\lceil t \rceil \leq t + 1$. $\square$

Corollary II.323.4 (Mass and one scalar functional). *Let an avoidance partition have R cells. For every nonnegative weight μ on H , some cell has mass at least $\sum_H \mu / R$. For every complex-valued function z on H , some cell satisfies*

$$\left| \sum_{x \in S} z(x) \right| \geq \frac{1}{R} \left| \sum_{x \in H} z(x) \right|.$$

Proof. The first assertion is averaging, and the second is the triangle inequality over the partition. $\square$

The favorable cell may depend on the weight or scalar functional. The corollary does not assert simultaneous preservation of an arbitrary list of unrelated correlations.

II.324 Small differences and a vertical-defect budget

Source 61, delivered Section 4.

We include the standard growth ingredients to make the exponents in the restriction theorem independently checkable. The minimal-growth argument is due to Petridis [36]; it is used here as a known theorem with a reproduced proof, not as a new ingredient.

Lemma II.324.1 (Minimal-growth lemma). *Let A, B be finite nonempty subsets of an abelian group. Choose a nonempty $X \subseteq A$ minimizing $|X + B|/|X|$, and write $\kappa = |X + B|/|X|$. Then for every finite C ,*

$$|X + B + C| \leq \kappa |X + C|.$$

Proof. The empty case is immediate. Otherwise order $C = \{c_1, \dots, c_t\}$, and define

$$X_i = \{x \in X : x + c_i \notin X + \{c_1, \dots, c_{i-1}\}\}.$$

The newly added parts $X_i + c_i$ are disjoint and cover $X + C$, so $|X + C| = \sum_i |X_i|$.

Every point of $(X \setminus X_i) + B + c_i$ is already in $X + B + \{c_1, \dots, c_{i-1}\}$. Indeed, an equality $x + c_i = x' + c_j$ with $j < i$ remains an equality after addition of the same $b \in B$. Thus the number of new points contributed by $X + B + c_i$ is at most

$$|X + B| - |(X \setminus X_i) + B|.$$

By minimality, $|(X \setminus X_i) + B| \geq \kappa |X \setminus X_i|$, with the empty case interpreted as zero. The number of new points is therefore at most $\kappa |X_i|$. Summing over i proves the lemma. $\square$

Lemma II.324.2 (Ruzsa triangle inequality). *If U, V, X are finite subsets of an abelian group and $X \neq \emptyset$, then*

$$|U - V| |X| \leq |U + X| |V + X|.$$

Proof. For each $d \in U - V$ choose one representation $d = u_d - v_d$. The map

$$(d, x) \mapsto (u_d + x, v_d + x)$$

is injective into $(U + X) \times (V + X)$. The difference of the two coordinates recovers d , and then either coordinate recovers x . $\square$

Proposition II.324.3 (Plünnecke–Ruzsa difference estimate). *Let A be finite nonempty and $|A - A| \leq K|A|$. Then for all nonnegative integers r, s ,*

$$|rA - sA| \leq K^{r+s} |A|. \quad (\text{II.324.1})$$

Proof. Apply Lemma II.324.1 with $B = -A$ and a minimizing nonempty $X \subseteq A$. Its growth ratio satisfies $\kappa \leq K$. Repeated application, with $C = (j - 1)B$, gives

$$|X + jB| \leq \kappa^j |X| \quad (j \geq 0).$$

Lemma II.324.2 now gives

$$|rB - sB| \leq \frac{|X + rB| |X + sB|}{|X|} \leq \kappa^{r+s} |X| \leq K^{r+s} |A|.$$

Negation identifies the cardinality on the left with $|rA - sA|$. $\square$

Lemma II.324.4 (Disjoint graph translates). *Let Γ be the graph of $\phi : B \rightarrow H$, where $B \subseteq G$ is nonempty. If $|\Gamma - \Gamma| \leq K|B|$, then*

$$|\mathcal{D}_m(\Gamma)| \leq K^{2m+1}. \quad (\text{II.324.2})$$

More precisely,

$$|B| |\mathcal{D}_m(\Gamma)| \leq |(m + 1)\Gamma - m\Gamma|.$$

Proof. For different $h \in H$, the vertical translates $\Gamma + (0, h)$ are pairwise disjoint. If $(x, \phi(x) + h) = (y, \phi(y) + h')$, the first coordinate gives $x = y$, and then the second gives $h = h'$. For $h \in \mathcal{D}_m(\Gamma)$ each translate lies in $(m + 1)\Gamma - m\Gamma$. Counting their disjoint union proves the precise statement; Proposition II.324.3 gives (II.324.2). $\square$

This packing step uses the fact that Γ is a graph. Applying only $|m\Gamma - m\Gamma| \leq K^{2m} |\Gamma|$ would leave an unwanted factor $|B|$ in the defect budget. The extra translate removes that dependence at the cost of one power of K .

II.325 Exact Freiman partitions and retention statements

Source 61, delivered Section 5.

Proof of Theorem F.66.1. The defect set $\mathcal{D}_m(\Gamma)$ contains zero and is symmetric. Apply Theorem II.323.3 in H to $D = \mathcal{D}_m(\Gamma) \setminus \{0\}$, and pull its color classes back by ϕ . Discard empty inverse images. The resulting partition of B uses at most $(5m)^{\lceil \log_2 |\mathcal{D}_m(\Gamma)| \rceil}$ cells.

Suppose $x_1, \dots, x_m, y_1, \dots, y_m$ lie in a cell B_j and have equal base sums. Their image defect

$$h = \sum_i \phi(x_i) - \sum_i \phi(y_i)$$

belongs to $\mathcal{D}_m(\Gamma)$ by definition. All the image points lie in one avoidance cell, so $h \notin D$. Hence $h = 0$, proving exactness. Finally Lemma II.324.4 and (II.323.2) give $R \leq 5mK^{a_m}$. Averaging the cardinalities gives the size assertion. $\square$

Corollary II.325.1 (Balanced bounded-coefficient systems). *In each cell supplied by Theorem F.66.1, ϕ preserves every integer-coefficient equation*

$$\sum_{i=1}^t c_i x_i = 0, \quad \sum_i c_i = 0, \quad \sum_{c_i > 0} c_i \leq m.$$

It preserves any finite system of such equations on the same cell.

Proof. If all coefficients vanish, the assertion is immediate. Otherwise expand positive and negative coefficients as repetitions, obtaining an equal-sum relation of order at most m . Pad both sides equally to order m . Applying exactness gives the asserted image equation. The same cell is exact for every order- m relation, so the assertion holds for a whole system. $\square$

Corollary II.325.2 (Several maps with a joint hypothesis). *Let $\phi_i : B \rightarrow H_i$ for $1 \leq i \leq t$, and let $\Phi = (\phi_1, \dots, \phi_t)$. If the joint graph in $G \times \prod_i H_i$ has difference ratio at most K , the bound in Theorem F.66.1 supplies one partition on which every ϕ_i is exact of order m .*

Proof. Apply the theorem with codomain $\prod_i H_i$ and project the resulting image-sum equalities onto each coordinate. $\square$

There is no extra factor depending on t once the joint hypothesis is given. Separate small-difference or energy bounds for the individual graphs do not establish that hypothesis.

Proposition II.325.3 (Weighted energy retained by a partition). *Let $B = \bigsqcup_{j=1}^R B_j$ be a partition with $R \geq 1$, let $w : B \rightarrow \mathbb{C}$, and put $w_j = w \mathbf{1}_{B_j}$. Then, for every $s \geq 1$,*

$$\sum_{j=1}^R \mathcal{E}_s(w_j) \geq R^{1-2s} \mathcal{E}_s(w), \quad \max_j \mathcal{E}_s(w_j) \geq R^{-2s} \mathcal{E}_s(w). \quad (\text{II.325.1})$$

If the cells are exact restrictions of ϕ of order m , then for $s \leq m$ every energy on the left also equals its graph-weighted counterpart.

Proof. Define $\hat{w}(\chi) = \sum_{x \in G} w(x) \exp(-2\pi i \chi(x))$. Orthogonality and the convolution identity give

$$\mathcal{E}_s(w) = \frac{1}{|G|} \sum_{\chi \in \hat{G}} |\hat{w}(\chi)|^{2s}.$$

Minkowski's inequality for this normalized L^{2s} norm, followed by Hölder's inequality for the finite sum over cells, gives

$$\begin{aligned}\mathcal{E}_s(w)^{1/(2s)} &\leq \sum_j \mathcal{E}_s(w_j)^{1/(2s)} \\ &\leq R^{1-1/(2s)} \left(\sum_j \mathcal{E}_s(w_j) \right)^{1/(2s)}.\end{aligned}$$

Raise to the power $2s$, and then average to obtain the maximum bound. Weighted exactness is Lemma II.322.3. $\square$

Proposition II.325.4 (One correlation and a nonnegative budget). *For the partition in Theorem F.66.1, every function $z : B \rightarrow \mathbb{C}$ admits an exact cell satisfying*

$$\left| \sum_{x \in B_j} z(x) \right| \geq \frac{1}{5mK^{a_m}} \left| \sum_{x \in B} z(x) \right|.$$

Every nonnegative mass admits the analogous retention inequality.

Proof. Use the triangle inequality and averaging as in Corollary II.323.4, with $R \leq 5mK^{a_m}$. $\square$

For example, one can take $z(x) = f(x) \exp(-2\pi i \psi(x))$ for a specified test phase ψ . This preserves that single scalar sum on one exact restriction. It does not preserve all such tests at once, and it does not guarantee that an earlier energy-to-growth selection has retained the original test sum. These are separate issues.

II.325.1 An elementary lower energy bound

Lemma II.325.5 (Energy from sumset support). *For every nonempty finite A and $s \geq 1$,*

$$\mathcal{E}_s(A) \geq \frac{|A|^{2s}}{|sA|}.$$

Proof. The representation function $r_s = \mathbf{1}_A^{*s}$ is supported on sA and has total mass $|A|^s$. Cauchy–Schwarz gives $|A|^{2s} \leq |sA| \sum_z r_s(z)^2$. $\square$

This inequality will allow the same restriction to satisfy both the cardinality and energy conclusions in Theorem F.66.2. It avoids selecting different favorable cells for the two criteria.

II.326 A self-contained energy-to-growth input

Source 61, delivered Section 6.

The next proposition is a conservative version of a standard Balog–Szemerédi–Gowers step. Its proof uses the familiar popular-difference graph and common-neighborhood method. We include it so that no unverified numerical estimate in another research manuscript is needed for the principal constants. This is not an improvement over the strongest versions available in the inspected research collection [252].

Proposition II.326.1 (A conservative BSG estimate). *Let A be a nonempty finite subset of an abelian group, $n = |A|$, and $0 < c \leq 1$. If $\mathcal{E}_2(A) \geq cn^3$, there exists $A_1 \subseteq A$ such that*

$$|A_1| \geq cn/8, \quad |A_1 - A_1| \leq 2^{21}c^{-9}n. \quad (\text{II.326.1})$$

In particular, A_1 is nonempty and has difference ratio at most $2^{24}c^{-10}$.

Proof. Let

$$r(d) = \#\{(a, b) \in A^2 : a - b = d\}, \quad P = \{d : r(d) \geq cn/2\}, \quad \delta = c/2.$$

We have $\sum_d r(d) = n^2$ and $\sum_d r(d)^2 = \mathcal{E}_2(A)$. The contribution from $d \notin P$ is at most $(cn/2) \sum_d r(d) = cn^3/2$. Since $r(d) \leq n$,

$$\sum_{d \in P} r(d) \geq cn^2/2 = \delta n^2, \quad |P| \leq 2n/c = n/\delta. \quad (\text{II.326.2})$$

Form a bipartite graph whose two vertex classes are copies of A , joining a left vertex a to a right vertex b when $a - b \in P$. Write $N(a)$ for the right neighborhood of a left vertex. Call an ordered pair (a, a') of left vertices bad if

$$|N(a) \cap N(a')| < t, \quad t = \delta^2 n/32.$$

The pair may have equal entries; all counts remain ordered.

Choose a right vertex v uniformly and let U be its left neighborhood. By (II.326.2), $\mathbb{E}|U| \geq \delta n$. If $b(U)$ is the number of ordered bad pairs in U^2 , then

$$\mathbb{E}b(U) = \frac{1}{n} \sum_{(a, a') \text{ bad}} |N(a) \cap N(a')| \leq tn.$$

Therefore

$$\mathbb{E} \left(|U| - \frac{16b(U)}{\delta n} \right) \geq \delta n - \frac{16t}{\delta} = \frac{\delta n}{2}.$$

Fix v with score at least $\delta n/2$, and write $M = |U|$. It follows that

$$M \geq \delta n/2, \quad b(U) \leq \frac{\delta n M}{16} \leq \frac{M^2}{8}. \quad (\text{II.326.3})$$

Delete from U every vertex with more than $M/4$ bad partners in U . Since badness is symmetric and the pairs are ordered, summing the bad-partner counts gives exactly $b(U)$. At most $4b(U)/M \leq M/2$ vertices are deleted. Denote the remaining left vertices by A_1 . Then

$$|A_1| \geq M/2 \geq \delta n/4 = cn/8.$$

For any $u, v \in A_1$, at least $M/2$ vertices $w \in U$ are good partners of both endpoints. Each such w has at least t common right neighbors with u and at least t with v . Thus there are at least

$$\frac{Mt^2}{2} \geq \frac{\delta^5 n^3}{4096} \quad (\text{II.326.4})$$

length-four walks $u - b - w - b' - v$. Repeated vertices are allowed, as they are throughout this proof.

For each $d \in A_1 - A_1$, fix endpoints $u_d, v_d \in A_1$ with $u_d - v_d = d$. Encode a walk between these endpoints by

$$(p_1, p_2, p_3, p_4) = (u_d - b, w - b, w - b', v_d - b') \in P^4.$$

The label tuple determines the walk once its endpoints are fixed: recover b from u_d, p_1 , then w from b, p_2 , and then b' from w, p_3 . Moreover

$$p_1 - p_2 + p_3 - p_4 = d,$$

so label sets arising from different differences d are disjoint. Using (II.326.2) and (II.326.4),

$$|A_1 - A_1| \frac{\delta^5 n^3}{4096} \leq |P|^4 \leq (n/\delta)^4.$$

Rearranging gives $|A_1 - A_1| \leq 2^{12} \delta^{-9} n = 2^{21} c^{-9} n$. Dividing by $|A_1| \geq cn/8$ yields the final ratio bound. $\square$

No asymptotic qualification is hidden in this argument. The threshold t and the score may be nonintegral; comparisons with the integer counts are literal real inequalities. The argument remains valid when n is small and when the group has elements of order two.

Lemma II.326.2 (High energy controls ordinary additive energy). *For a finite set A of size n and every $m \geq 2$,*

$$\mathcal{E}_m(A) \leq n^{2m-4} \mathcal{E}_2(A). \quad (\text{II.326.5})$$

Proof. The inequality $\|u * v\|_2 \leq \|u\|_2 \|v\|_1$ follows by the triangle inequality applied to the sum of translates of u , weighted by v . Apply it repeatedly to $\mathbf{1}_A^{*m} = (\mathbf{1}_A * \mathbf{1}_A) * \mathbf{1}_A^{*(m-2)}$. Since $\|\mathbf{1}_A\|_1 = n$,

$$\|\mathbf{1}_A^{*m}\|_2 \leq n^{m-2} \|\mathbf{1}_A * \mathbf{1}_A\|_2.$$

Squaring proves the result, including equality of the two expressions when $m = 2$. $\square$

Thus the hypothesis $\mathcal{E}_m^\phi(B) \geq \alpha n^{2m-1}$ implies $\mathcal{E}_2(\Gamma) \geq \alpha n^3$. Applying Proposition II.326.1 to the graph, rather than to the domain alone, is crucial: the extracted small-difference set must still encode the map.

II.327 Polynomial exact restriction of high relation energy

Source 61, delivered Section 7.

Theorem II.327.1 (General-order energy restriction). *Let $m \geq 2$, $B \subseteq G$ be nonempty with $|B| = n$, and $\phi : B \rightarrow H$. Suppose $0 < \alpha \leq 1$ and either*

$$\mathcal{E}_m^\phi(B) \geq \alpha n^{2m-1} \quad \text{or} \quad \mathcal{E}_2^\phi(B) \geq \alpha n^3.$$

Put

$$a_m = (2m + 1) \log_2(5m), \quad \theta_m = 2ma_m + m, \quad K = 2^{24} \alpha^{-10}.$$

There is an exact order- m restriction $B' \subseteq B$ satisfying simultaneously

$$|B'| \geq \frac{\alpha n}{8(5m)K^{a_m}}, \quad (\text{II.327.1})$$

$$\mathcal{E}_m(B') = \mathcal{E}_m^\phi(B') \geq (5m)^{-2m} K^{-\theta_m} \left(\frac{\alpha n}{8} \right)^{2m-1}. \quad (\text{II.327.2})$$

Equivalently, the energy lower bound is

$$\mathcal{E}_m(B') \geq C_m \alpha^{p_m} n^{2m-1}, \quad \begin{cases} p_m = 2m - 1 + 10\theta_m, \\ C_m = (5m)^{-2m} 2^{-3(2m-1)-24\theta_m}. \end{cases} \quad (\text{II.327.3})$$

In particular, $p_m = O(m^2 \log m)$.

Proof. Lemma II.326.2 reduces the first hypothesis to the second. Proposition II.326.1, applied to Γ , gives a subset Γ_1 of size $M \geq \alpha n/8$ and difference ratio at most K . Because it is a subset of a graph, Γ_1 is the graph of ϕ restricted to a subset $B_1 \subseteq B$.

Partition B_1 by Theorem F.66.1, and choose a largest cell B' . Then $|B'| \geq M/R$, where $R \leq 5mK^{a_m}$. This proves (II.327.1). Projection onto G and Proposition II.324.3 give

$$|mB'| \leq |m\Gamma_1| \leq K^m M.$$

Lemma II.325.5 therefore yields

$$\begin{aligned} \mathcal{E}_m(B') &\geq \frac{(M/R)^{2m}}{K^m M} = R^{-2m} K^{-m} M^{2m-1} \\ &\geq (5m)^{-2m} K^{-(2ma_m+m)} \left(\frac{\alpha n}{8} \right)^{2m-1}. \end{aligned}$$

Every relation is respected by exactness. Substitution of $K = 2^{24}\alpha^{-10}$ gives (II.327.3). $\square$

II.327.1 A plug-in interface for stronger BSG inputs

Proposition II.327.2 (Modular transfer). *Let $m \geq 2$, $0 < \alpha \leq 1$, and use a_m, θ_m from Theorem II.327.1. Suppose an energy-to-growth theorem applicable to Γ supplies a subgraph Γ_1 with*

$$M = |\Gamma_1| \geq c_0 \alpha^u n, \quad |\Gamma_1 - \Gamma| \leq C \alpha^{-v} M,$$

where $c_0 > 0$, $u, v \geq 0$, and $C \geq 1$. Then an exact order- m restriction satisfies simultaneously

$$|B'| \geq \frac{c_0}{5mC^{a_m}} \alpha^{u+va_m} n, \quad (\text{II.327.4})$$

$$\mathcal{E}_m(B') \geq \frac{c_0^{2m-1}}{(5m)^{2m} C^{\theta_m}} \alpha^{u(2m-1)+v\theta_m} n^{2m-1}. \quad (\text{II.327.5})$$

Proof. Repeat the last proof with $K = C\alpha^{-v}$ and $M \geq c_0\alpha^u n$. Nothing else changes. $\square$

This is the correct place to import a stronger established energy-to-growth theorem. A statement bounding $|\Gamma_1 - \Gamma|$ relative to the *original* size n must first be converted to a bound relative to M . Losing this distinction changes both the constant and the exponent. The present unconditional numerical results use only Proposition II.326.1; the modular statement does not certify the hypotheses of another source.

II.327.2 The order-eight calculation

Proof of Theorem F.66.2. Use the preceding proof with $m = 8$. The exact integer inequality

$$40^{17} < 2^{91}$$

gives $a_8 = 17 \log_2 40 < 91$. Consequently

$$R \leq 40K^{a_8} \leq 2^6 K^{91}, \quad M \geq \alpha n/8, \quad K = 2^{24} \alpha^{-10}.$$

The largest cell therefore has

$$|B'| \geq \frac{M}{R} \geq 2^{-9} \alpha n K^{-91} = 2^{-(9+24 \cdot 91)} \alpha^{1+10 \cdot 91} n = 2^{-2193} \alpha^{911} n.$$

The same cell has energy at least

$$\begin{aligned} R^{-16} K^{-8} M^{15} &\geq 2^{-96} K^{-(16 \cdot 91 + 8)} (\alpha n/8)^{15} \\ &= 2^{-141} K^{-1464} \alpha^{15} n^{15} \\ &= 2^{-(141 + 24 \cdot 1464)} \alpha^{15 + 10 \cdot 1464} n^{15} \\ &= 2^{-35277} \alpha^{14655} n^{15}. \end{aligned}$$

Exactness gives equality of the base and respected energy counts. The proof starts equally well from either hypothesis in Theorem F.66.2. $\square$

Quantity	Value used in the order-eight proof
Graph-core size	$M \geq \alpha n/8$
Graph-core difference ratio	$K \leq 2^{24} \alpha^{-10}$
Partition size	$R \leq 2^6 K^{91}$
Sumset support	$ 8B' \leq K^8 M$
Size exponents	$9 + 24 \cdot 91 = 2193, 1 + 10 \cdot 91 = 911$
Energy exponents	$141 + 24 \cdot 1464 = 35277$ $15 + 10 \cdot 1464 = 14655$

II.328 Consequences for Gowers's restriction lemmas

Source 61, delivered Section 8.

II.328.1 Lemma 7.5: a generalization, not a prime-field constant improvement

Gowers's Lemma 7.5 extracts an order- m Freiman restriction from a graph with small difference set. The original proof uses scalar multiplication and an interval in the codomain. The inspected formal statement explicitly requires the modulus to be prime [1, 250].

Theorem F.66.1 replaces the prime scalar by a short list of characters and replaces the single selected interval by a complete partition. Its codomain can be, for example,

$$H = (\mathbb{Z}/8\mathbb{Z})^{r_1} \times (\mathbb{Z}/9\mathbb{Z})^{r_2} \times (\mathbb{Z}/25\mathbb{Z})^{r_3},$$

with no explicit rank or exponent factor. The domain need not be cyclic either.

The cost is visible: the exponent $a_m = (2m+1)\log_2(5m)$ is larger than $4m$ for the relevant values of m . The uniform theorem is not advertised as a numerical improvement in the prime-modulus case. Its output is stronger in generality and in the partition structure, and its precise defect-budget version can be much better than its worst-case K bound.

II.328.2 Lemma 9.3: exactness and removal of the threshold

Corollary II.328.1 (A strengthened cyclic Lemma 9.3). *Let $N \geq 1$, $B \subseteq \mathbb{Z}/N\mathbb{Z}$ be nonempty, $|B| = \beta N$, and $\phi : B \rightarrow \mathbb{Z}/N\mathbb{Z}$. Suppose*

$$\mathcal{E}_8^\phi(B) \geq \alpha \beta^{15} N^{15}, \quad \alpha > 0.$$

Then there is an exact order-eight restriction $B' \subseteq B$ with

$$\mathcal{E}_8(B') \geq 2^{-35277} \alpha^{14655} \beta^{15} N^{15}.$$

For every $0 < \eta \leq 1$, this implies the numerical and approximate-respect conclusions of the inspected statement of Lemma 9.3, with $N_0 = 1$.

Proof. The energy cap (II.322.1) implies $\alpha \leq 1$. Apply Theorem F.66.2 with $n = \beta N$. Exactness implies any requirement that at least a proportion $1 - \eta$ of the additive tuples be respected. Finally,

$$\begin{aligned} (\alpha\eta/4)^{2^{19}} &= 2^{-1048576} \alpha^{524288} \eta^{524288} \\ &\leq 2^{-35277} \alpha^{14655}, \end{aligned}$$

since $0 < \alpha, \eta \leq 1$, $1048576 \geq 35277$, and $524288 \geq 14655$. No restriction on N beyond positivity was used. $\square$

[write] **7 October 2026; sources 60–69.** *Against the formal statement* (checked at the write). For nonempty B the hypothesis forces $\alpha \leq 1$, and for $0 < \alpha, \eta \leq 1$ one has $(\alpha\eta/4)^{2^{19}} \leq 2^{-2^{20}} \alpha^{2^{19}} \leq 2^{-35277} \alpha^{14655}$, since $2^{20} = 1,048,576 \geq 35,277$ and $2^{19} = 524,288 \geq 14,655$; exactness gives $\text{GammaHomOfOrder } 8 \ B' \ \phi(1 - \eta)$. So this corollary implies `lemma_9_3` (Sections08_09.lean, line 77) with $N_0 = 1$, on paper; the catalogue statement is already exact (`lemma_9_3_holds`), and this corollary is not formalized.

The conclusion is stronger in three separate senses: its displayed coefficient is larger throughout the stated parameter range, it asserts exact order-eight respect, and it has no large-modulus premise. The proof does not assert that its coefficient dominates the sharper *prime-field* bounds in other project reports. Those comparisons involve different hypotheses and different outputs.

II.328.3 Corollary 9.4: no auxiliary density penalty

Corollary II.328.2 (Direct energy version of Corollary 9.4). *Let G, H be arbitrary finite abelian groups and $B \subseteq G$ nonempty. If*

$$\mathcal{E}_2^\phi(B) \geq \gamma |B|^3, \quad \gamma > 0,$$

then there is an exact order-eight restriction $B' \subseteq B$ with

$$\mathcal{E}_8(B') \geq 2^{-35277} \gamma^{14655} |B|^{15}.$$

Proof. The energy cap implies $\gamma \leq 1$. Use the second hypothesis allowed in Theorem F.66.2. $\square$

[write] **7 October 2026; sources 60–69.** *Against the formal statement* (checked at the write). `corollary_9_4` (Sections08_09.lean, line 91) assumes `GammaAdditive B` $\phi(\gamma\beta^3)$, that is at least $\gamma|B|^3$ respected quadruples, so $\gamma \leq 1$, and asks for $(\gamma^7\beta^6\eta/4)^{2^{19}}\beta^{15}N^{15}$ additive sixteen-tuples on an approximate homomorphism; since $(\gamma^7\beta^6\eta/4)^{2^{19}} \leq 2^{-2^{20}}\gamma^{7 \cdot 2^{19}} \leq 2^{-35277}\gamma^{14655}$, this corollary implies it with $N_0 = 1$, on paper; it is already exact (`corollary_9_4_holds`, Proofs09Restriction.lean, line 416).

For $G = \mathbb{Z}/N\mathbb{Z}$ and $|B| = \beta N$, this is $2^{-35277}\gamma^{14655}\beta^{15}N^{15}$. The inspected Corollary 9.4 instead passes through a normalized sixteen-tuple hypothesis and displays

$$(\gamma^7\beta^6\eta/4)^{2^{19}}\beta^{15}N^{15}$$

with a size threshold [250]. The direct graph-growth route avoids both the auxiliary β^6 penalty inside the large power and the approximate-respect parameter. Again, it does not by itself preserve the later analytic data associated with the graph.

II.328.4 Why the original density normalization survives

If one knew only $|B'| \geq \rho|B|$, the elementary ambient-group bound would give

$$\mathcal{E}_8(B') \geq \frac{|B'|^{16}}{|G|},$$

which in a cyclic group is of scale $\rho^{16}\beta^{16}N^{15}$. This has an extra factor of β compared with the normalization in Lemma 9.3. The proof above instead uses $|8B'| \leq K^8M$, where M is the size of the selected *graph core*. It therefore produces a multiple of M^{15} and, ultimately, of $|B|^{15}$. This support estimate is the quantitative reason the restriction can retain the desired relation-energy scale.

II.329 Torsion obstructions and the optimal logarithmic scale

Source 61, delivered Section 9.

The absence of an ambient-order parameter in Theorem II.323.3 does not mean that torsion is harmless. Rather, torsion is paid for by the exponent of the defect budget. The following examples show that logarithmic growth of that exponent in m is necessary.

II.329.1 A coloring problem requiring singleton cells

Let $\chi_m(H, D)$ denote the least number of cells in a partition of H satisfying $(mS - mS) \cap D = \emptyset$ for every cell S .

Proposition II.329.1 (An exact 2-primary obstruction). *Let $e \geq 1$, $r \geq 1$, and $2^{e-1} \leq m$. Set*

$$H = (\mathbb{Z}/2^e\mathbb{Z})^r, \quad T = H[2] = 2^{e-1}H, \quad D = T \setminus \{0\}.$$

Then every admissible cell is a singleton, and

$$\chi_m(H, D) = 2^{er} = (|D| + 1)^e. \tag{II.329.1}$$

Proof. Let $x \neq y$. Their difference has order 2^j for some $1 \leq j \leq e$. Put $q = 2^{j-1} \leq m$. Then $q(x - y)$ is a nonzero element of $H[2]$, hence belongs to D .

If a cell contains x, y , the difference $q(x - y)$ belongs to $mS - mS$: use q copies of x against q copies of y , and append the same $m - q$ copies of y to both sides. Thus any two distinct elements in one cell violate avoidance. Singleton cells are admissible because $m\{x\} - m\{x\} = \{0\}$. Their number is $|H| = 2^{er}$, while $|D| + 1 = 2^r$. $\square$

Proof of Theorem F.66.3. Theorem II.323.3 gives the upper bound with constant $C_m = 5m$. For the lower bound choose $e = \lfloor \log_2 m \rfloor + 1$, so that $2^{e-1} \leq m$, and let r tend to infinity in Proposition II.329.1. If a universal bound with exponent $b < e$ held, it would require

$$2^{er} \leq C_m 2^{br} \quad \text{for every } r,$$

which is impossible. The limiting statements follow because the two bounds differ from $\log_2 m$ by bounded amounts. $\square$

Remark II.329.2 (The case $m = 1$). For symmetric D , the ordinary Cayley graph with differences in D has maximum degree $|D|$. A greedy vertex coloring uses at most $|D| + 1$ colors. Taking $D = H \setminus \{0\}$ forces singleton cells. Thus the analogous exponent is exactly $b_1 = 1$. The character bound remains valid but is not sharp in this case.

Proposition II.329.3 (A p -primary variant). *Let p be prime, $p^{e-1} \leq m$, $H = (\mathbb{Z}/p^e\mathbb{Z})^r$, and $D = H[p] \setminus \{0\}$. Then*

$$\chi_m(H, D) = p^{er} = (|D| + 1)^e.$$

Proof. For $x \neq y$ with difference of order p^j , multiply the difference by $p^{j-1} \leq m$. It becomes a nonzero element of $H[p]$. The padding argument and the singleton upper bound are identical to those above. $\square$

The prime 2 yields the largest lower exponent from this family for a given m . The obstruction uses increasing rank and fixed exponent; it cannot be removed simply by checking sufficiently large moduli.

II.329.2 Graph restrictions that must be singletons

The preceding coloring example allows an arbitrary forbidden set. The next construction shows that related torsion obstructions actually arise from graphs, and quantifies their graph difference ratio and energy.

Theorem II.329.4 (Canonical-section graph obstruction). *Let $e \geq 2$, $L = 2^{e-1} \leq m$, and $r \geq 1$. Define*

$$G = (\mathbb{Z}/L\mathbb{Z})^r, \quad H = (\mathbb{Z}/(2L)\mathbb{Z})^r,$$

and let $\phi : G \rightarrow H$ take each coordinate to its representative in $\{0, \dots, L-1\}$. With Γ its graph and $n = L^r$,

$$|\Gamma - \Gamma| = (2L - 1)^r, \quad K = (2 - 1/L)^r, \quad (\text{II.329.2})$$

$$\mathcal{E}_2(\Gamma) = \left(\frac{2L^3 + L}{3} \right)^r, \quad \frac{\mathcal{E}_2(\Gamma)}{n^3} = \left(\frac{2}{3} + \frac{1}{3L^2} \right)^r. \quad (\text{II.329.3})$$

Every exact Freiman restriction of order m has at most one point.

Proof. First consider one coordinate. The graph differences are precisely

$$(t \bmod L, t \bmod 2L), \quad -(L-1) \leq t \leq L-1.$$

These $2L-1$ pairs are distinct. Taking independent coordinates proves (II.329.2).

For two-term sums, all representative sums lie between 0 and $2L-2$. Equality modulo $2L$ is therefore equality as integers, which also implies equality modulo L . The one-coordinate graph energy is the additive energy of the interval $\{0, \dots, L-1\}$:

$$L^2 + 2 \sum_{j=1}^{L-1} j^2 = L^2 + \frac{(L-1)L(2L-1)}{3} = \frac{2L^3 + L}{3}.$$

Multiplication over the coordinates proves (II.329.3).

Now let $x \neq y$ in G and put $d = \phi(x) - \phi(y) \in H$. This difference is nonzero. Its order is 2^j for some $j \leq e$. In fact $j \geq 2$: every nonzero element of order two has some coordinate equal to L modulo $2L$, whereas each representative coordinate difference lies between $-(L-1)$ and $L-1$.

Set $q = 2^{j-1} \leq L \leq m$. Then qd is a nonzero element of $H[2]$, and $H[2]$ is the kernel of reduction from H to G . Consequently

$$q(x-y) = 0 \quad \text{in } G, \quad q(\phi(x) - \phi(y)) \neq 0 \quad \text{in } H.$$

This is an unrespected order- q relation using only x, y . Padding gives an unrespected order- m relation on the same pair. No exact restriction can contain both points. $\square$

Corollary II.329.5 (Necessary exponents for graph extraction). *Fix $m \geq 2$ and a power of two $L \leq m$, $L \geq 2$. If a universal theorem guarantees an exact order- m restriction of size at least $c_m K^{-a} |B|$ whenever the graph difference ratio is at most K , with $c_m > 0$ independent of the groups, then necessarily*

$$a \geq \frac{\log L}{\log(2 - 1/L)}. \quad (\text{II.329.4})$$

If a universal theorem instead guarantees size at least $c_m \gamma^u |B|$ under $\mathcal{E}_2^\phi(B) \geq \gamma |B|^3$, then necessarily

$$u \geq \frac{\log L}{-\log(2/3 + 1/(3L^2))}. \quad (\text{II.329.5})$$

Proof. Apply the proposed guarantee to Theorem II.329.4, and let $r \rightarrow \infty$. Since every exact restriction has size at most one, the bases in

$$c_m [L(2 - 1/L)^{-a}]^r \leq 1, \quad c_m [L(2/3 + 1/(3L^2))^u]^r \leq 1$$

must be at most one. Taking logarithms gives the conclusions. $\square$

Choosing the largest power of two below m makes both necessary exponents of order $\log m$. For the K exponent, the leading scale is $\log_2 m$. The proven upper graph exponent is $a_m = O(m \log m)$, leaving a factor of order m unresolved. The matching leading logarithm in Theorem F.66.3 belongs to the *defect-coloring* problem, not to this stronger graph-extraction optimization.

Example II.329.6 (The order-eight section). For $m = 8$, take $L = 8$. Then $G = (\mathbb{Z}/8\mathbb{Z})^r$, $H = (\mathbb{Z}/16\mathbb{Z})^r$, $n = 8^r$,

$$K = (15/8)^r, \quad \gamma = (43/64)^r.$$

Every exact order-eight restriction is a singleton. This gives a concrete test family for any proposed improvement to the order-eight size exponent. It is not a counterexample to Theorem F.66.2; the lower bounds there are compatible with singleton output on this family.

II.330 A translated Fejér restriction theorem

Source 60, delivered Section 9.

The restriction step in Section 9 of Gowers's proof turns a map that respects many additive relations into one that respects nearly all relations on a smaller set [1]. Source 04 replaces repeated elementary filters by one positive Fourier kernel and retains the Fourier weights of unwanted relations [244]. Source 17 uses a Fejér kernel whose amplification grows linearly with its bandwidth [245]. We combine those two features with an independent translation of the phase. The translation reduces the *total weight* of unwanted coefficient relations by a factor of the bandwidth. The same factor controls the survival of repeated tuples, so it is not lost when the selected set is interpreted as an actual subset.

The resulting prime-cyclic theorem has the retention coefficient of source 17, with a size threshold of degree $m - 1$ in the inverse accuracy parameter for m -term relations. The proposed composition of sources 04 and 17 had degree m . We also prove a finite-data certificate on arbitrary finite abelian groups, with small-order torsion and repeated tuples explicitly accounted for. The improvement is local to the restriction step; no later inverse-theorem parameter comparison is assumed.

II.330.1 Relations, the kernel, and the finite probability space

Let G, H be finite abelian groups, let $\emptyset \neq B \subseteq G$, write $b = |B|$, and let $\phi : B \rightarrow H$. Fix an even integer $m = 2r \geq 4$ and put

$$\varepsilon = (\underbrace{1, \dots, 1}_r, \underbrace{-1, \dots, -1}_r).$$

An ordered tuple $\mathbf{x} = (x_1, \dots, x_m) \in B^m$ is *additive* if $\sum_i \varepsilon_i x_i = 0$ in G , and is *respected* if also

$$\Delta_\phi(\mathbf{x}) := \sum_i \varepsilon_i \phi(x_i) = 0 \quad \text{in } H.$$

Let $T_m(B)$ count all additive tuples and let $G_m(B, \phi)$ count the respected ones. Repeated coordinates are included in both counts. For the current input abbreviate

$$T = T_m(B), \quad G_0 = G_m(B, \phi), \quad U = T - G_0,$$

and let U_{rep} count the unrespected additive tuples with at least one repeated coordinate. The elementary bounds

$$T \leq b^{m-1}, \quad U_{\text{rep}} \leq \binom{m}{2} b^{m-2} \tag{II.330.1}$$

hold on every finite abelian group. For the first, the last coordinate is uniquely determined by the preceding ones. For the second, prescribe an equality $x_i = x_j$ and choose a position $k \notin \{i, j\}$. Fixing the $m - 2$ positions outside $\{j, k\}$ determines x_j by the equality and x_k by its coefficient $\varepsilon_k = \pm 1$ in the additive relation. A union bound over the pairs proves the assertion, without any torsion hypothesis.

Write $\mathbb{T} = \mathbb{R}/\mathbb{Z}$ and $e(t) = \exp(2\pi it)$. For an integer $L \geq 2$ define

$$K_L(t) = \frac{1}{L^2} \left| \sum_{j=0}^{L-1} e(jt) \right|^2 = \sum_{n=1-L}^{L-1} a_n e(nt), \quad a_n = \frac{L - |n|}{L^2}. \tag{II.330.2}$$

Then $0 \leq K_L \leq 1$, the coefficients are nonnegative, $\sum_n a_n = 1$, and $\int_{\mathbb{T}} K_L = 1/L$. Set

$$a = L^{-m},$$

$$S = S_{m,L} = \sum_{|t| < L} \left(1 - \frac{|t|}{L}\right)^m, \quad (\text{II.330.3})$$

$$J = J_{m,L} = \int_{\mathbb{T}} K_L(t)^m dt. \quad (\text{II.330.4})$$

The useful bounds are

$$\frac{2L}{m+1} \leq S \leq L, \quad aS \leq J \leq \frac{1}{L}. \quad (\text{II.330.5})$$

The lower bound for S is the trapezoidal inequality for the convex function $(1 - t/L)^m$ on $[0, L]$; its upper bound follows by replacing the m th power by the first power. The inequality $J \leq 1/L$ follows from $K_L^m \leq K_L$. In the nonnegative Fourier expansion for J , the coefficient vectors $n = t\varepsilon$ all have coordinate sum zero and have total weight aS , proving the remaining inequality.

An exact formula, useful when optimizing finite parameters, is

$$J_{m,L} = L^{-2m} \sum_{j=0}^m (-1)^j \binom{2m}{j} \binom{(m-j)L + m - 1}{2m - 1}. \quad (\text{II.330.6})$$

Here a binomial coefficient is zero when its nonnegative upper argument is smaller than its lower argument. Indeed $L^{2m} J_{m,L}$ is the coefficient of $z^{m(L-1)}$ in $(1 + z + \cdots + z^{L-1})^{2m}$; expanding $(1 - z^L)^{2m}(1 - z)^{-2m}$ gives (II.330.6).

For a positive integer q define the leakage signal

$$S_{m,L}[q] = \sum_{\substack{|t| < L \\ q|t}} \left(1 - \frac{|t|}{L}\right)^m. \quad (\text{II.330.7})$$

In particular $S_{m,L}[q] = 1$ when $q > L - 1$. For a finite abelian group put

$$G[d] = \{x \in G : dx = 0\}, \quad \tau_G(L) = \max_{1 \leq d \leq 2L-2} |G[d]|. \quad (\text{II.330.8})$$

For $G = \mathbb{Z}/n\mathbb{Z}$ this is $\max_{1 \leq d \leq 2L-2} \gcd(n, d)$. It is 1 whenever multiplication by every integer between 1 and $2L - 2$ is injective on G .

All random choices can be finite. Put

$$Q = m(L - 1) + 1. \quad (\text{II.330.9})$$

Choose independently a uniform character $\chi \in \widehat{G}$, a uniform character $\psi \in \widehat{H}$, and a uniform integer $j \in \{0, \dots, Q - 1\}$. We regard characters additively as maps to $\mathbb{T}$. Conditional on these choices, include each point $x \in B$ independently in a random subset B' with probability

$$p(x) = K_L(\chi(x) + \psi(\phi(x)) + j/Q). \quad (\text{II.330.10})$$

The bound $0 \leq K_L \leq 1$ makes these valid Bernoulli probabilities. Let X and Y be the numbers of respected and unrespected additive m -tuples in B' , respectively.

II.330.2 An exact finite-group survival estimate

Theorem II.330.1 (Translated-kernel survival estimates). *With the definitions above, on arbitrary finite abelian groups G, H ,*

$$\mathbb{E}X \geq aSG_0, \quad (\text{II.330.11})$$

$$\begin{aligned} \mathbb{E}Y \leq a \sum_{\substack{\mathbf{x} \in B^m \text{ additive} \\ \Delta_\phi(\mathbf{x}) \neq 0}} S_{m,L}[\text{ord}(\Delta_\phi(\mathbf{x}))] \\ + \tau_G(L)(J - aS)b^{m-2} + (L^{-1} - a)U_{\text{rep}}. \end{aligned} \quad (\text{II.330.12})$$

In particular, if every nonzero discrepancy appearing among the unrespected tuples has order greater than $L - 1$, then

$$\mathbb{E}Y \leq aU + \frac{\tau_G(L) + \binom{m}{2}}{L} b^{m-2}. \quad (\text{II.330.13})$$

Proof. For a fixed tuple $\mathbf{x}$, write

$$P(\mathbf{x}) = \prod_{i=1}^m p(x_i), \quad R(\mathbf{x}) = \prod_{z \in \{x_1, \dots, x_m\}} p(z).$$

The first product includes multiplicities. The second is the true conditional probability that every point of the tuple is selected. Thus $R \geq P$, and equality holds when the coordinates are distinct.

For $n = (n_1, \dots, n_m) \in \{1 - L, \dots, L - 1\}^m$, put $w_n = \prod_i a_{n_i}$. Expanding the first product and applying orthogonality of the two finite dual groups and of the Q shifts gives

$$\mathbb{E}P(\mathbf{x}) = \sum_n w_n 1_{\{\sum_i n_i = 0 \text{ in } \mathbb{Z}\}} 1_{\{\sum_i n_i x_i = 0 \text{ in } G\}} 1_{\{\sum_i n_i \phi(x_i) = 0 \text{ in } H\}}. \quad (\text{II.330.14})$$

To justify the integer condition, the shift average initially requires $Q \mid \sum_i n_i$, whereas $|\sum_i n_i| \leq m(L - 1) < Q$. Hence divisibility is equivalent to exact integer equality. In particular, there is no assumption here that the group exponents exceed the kernel bandwidth.

Call $n = t\varepsilon$ an integer principal-line vector. On every respected additive tuple all these vectors contribute to (II.330.14), and their total weight is aS . The inequality $R \geq P$ therefore proves (II.330.11), including for tuples with repeated coordinates. On an unrespected tuple a principal-line vector contributes exactly when $t\Delta_\phi(\mathbf{x}) = 0$. Its total contribution is $aS_{m,L}[\text{ord}(\Delta_\phi(\mathbf{x}))]$.

Suppose next that n is not on the integer principal line. Some 2-by-2 minor of the two integer rows ε, n is then a nonzero integer d with $|d| \leq 2L - 2$. After fixing the other $m - 2$ coordinates in B , elimination using a coefficient $\varepsilon_i = \pm 1$ reduces the equations

$$\sum_i \varepsilon_i x_i = 0, \quad \sum_i n_i x_i = 0$$

to an equation $dz = z_0$ for one remaining coordinate. Such an equation has either no solutions or $|G[d]|$ solutions, with the other coordinate uniquely determined. Thus there are at most $\tau_G(L)b^{m-2}$ additive tuples satisfying the additional domain relation. Imposing the image relation, or restricting to unrespected tuples, cannot increase this count.

The total weight of all vectors with $\sum_i n_i = 0$ is J . Subtracting the principal-line weight shows that the nonprincipal vectors have total weight $J - aS$. Summing first over tuples and then over these vectors consequently bounds their total contribution by $\tau_G(L)(J - aS)b^{m-2}$.

It remains to compare multiplicity products with actual survival. For any chosen coordinate x_i one has $R(\mathbf{x}) \leq p(x_i)$. For fixed χ, ψ , averaging over the Q shifts gives $\mathbb{E}_j p(x_i) = 1/L$, since $Q > L - 1$. Thus $\mathbb{E}R(\mathbf{x}) \leq 1/L$. The zero coefficient vector in (II.330.14) contributes a on every tuple, so $\mathbb{E}P(\mathbf{x}) \geq a$. Therefore

$$0 \leq \mathbb{E}(R(\mathbf{x}) - P(\mathbf{x})) \leq L^{-1} - a.$$

This difference is zero for distinct-coordinate tuples. Summing it over the U_{rep} unrespected repeated tuples proves (II.330.12). Finally use (II.330.1), $J - aS \leq 1/L$, and $S_{m,L}[q] = 1$ for $q > L - 1$ to obtain (II.330.13). $\square$

Remark II.330.2 (A finer domain torsion statistic). The factor $\tau_G(L)(J - aS)$ in (II.330.12) can be replaced by

$$\sum_{\substack{n \notin \mathbb{Z}\varepsilon \\ \sum_i n_i = 0}} w_n \min_{\substack{i < j \\ d_{ij} \neq 0}} |G[d_{ij}]|, \quad d_{ij} = \varepsilon_i n_j - \varepsilon_j n_i.$$

The proof simply chooses the best nonzero minor separately for each coefficient vector. The actual image-discrepancy orders and U_{rep} can likewise be retained. Thus the theorem gives a certificate sensitive to the input's torsion and relation statistics, rather than requiring every one of them to be replaced by a uniform maximum.

II.330.3 Restriction with a simultaneous count guarantee

Corollary II.330.3 (Finite-data certificate). *Let $0 < \eta \leq 1$ and define*

$$\begin{aligned} \Psi_L = \eta a S G_0 - a \sum_{\substack{\mathbf{x} \in B^m \text{ additive} \\ \Delta_\phi(\mathbf{x}) \neq 0}} S_{m,L}[\text{ord}(\Delta_\phi(\mathbf{x}))] \\ - \tau_G(L)(J - aS)b^{m-2} - (L^{-1} - a)U_{\text{rep}}. \end{aligned}$$

If $\Psi_L > 0$, some $B' \subseteq B$ satisfies

$$T_m(B') \geq \frac{\Psi_L}{\eta} > 0, \quad \frac{G_m(B', \phi)}{T_m(B')} \geq \frac{1}{1 + \eta} \geq 1 - \eta. \quad (\text{II.330.15})$$

Proof. Theorem II.330.1 gives $\mathbb{E}(\eta X - Y) \geq \Psi_L$. Choose a realization with $\eta X - Y \geq \Psi_L$. Then $Y \leq \eta X$ and

$$T_m(B') = X + Y \geq X \geq \frac{\eta X - Y}{\eta} \geq \frac{\Psi_L}{\eta}.$$

These inequalities prove the simultaneous proportion and count claims. $\square$

Theorem II.330.4 (An explicit polynomial restriction). *Suppose $0 < \alpha, \eta \leq 1$ and $G_m(B, \phi) \geq \alpha b^{m-1}$. Put*

$$L = \left\lceil \frac{m+1}{\alpha\eta} \right\rceil, \quad C = \tau_G(L) + \binom{m}{2}.$$

Assume that every nonzero discrepancy on an unrespected additive tuple has order greater than $L - 1$, and that

$$b \geq 2CL^{m-1}. \quad (\text{II.330.16})$$

Then there is a subset $B' \subseteq B$ such that

$$\frac{G_m(B', \phi)}{T_m(B')} \geq \frac{1}{1 + \eta} \geq 1 - \eta, \quad (\text{II.330.17})$$

$$T_m(B') \geq \frac{b^{m-1}}{2\eta L^m} \geq \frac{\alpha^m \eta^{m-1}}{2(m+2)^m} b^{m-1}. \quad (\text{II.330.18})$$

Proof. By (II.330.5), $\eta\alpha S \geq 2$. The elementary upper bound for T and (II.330.13) give

$$\begin{aligned} \mathbb{E}(\eta X - Y) &\geq a(\eta\alpha S - 1)b^{m-1} - \frac{C}{L}b^{m-2} \\ &\geq ab^{m-1} - \frac{C}{L}b^{m-2} \geq \frac{a}{2}b^{m-1}. \end{aligned}$$

The last inequality is exactly (II.330.16). The signed-objective argument of Corollary II.330.3 proves the first bound in (II.330.18) and the proportion claim. Finally $L \leq (m+2)/(\alpha\eta)$ gives the second retention bound. $\square$

Corollary II.330.5 (Prime cyclic groups). *Let $B \subseteq \mathbb{F}_p$ and $\phi : B \rightarrow \mathbb{F}_p$, where p is prime. Under the hypothesis $G_m(B, \phi) \geq \alpha b^{m-1}$, define L as in Theorem II.330.4. The sole size condition*

$$b \geq 2 \left(1 + \binom{m}{2} \right) L^{m-1} \quad (\text{II.330.19})$$

implies the conclusions (II.330.17)–(II.330.18).

Proof. The condition implies $p \geq b > 2L - 2$. Consequently $\tau_G(L) = 1$ and every nonzero image discrepancy has order $p > L - 1$. Apply Theorem II.330.4. $\square$

[write] **7 October 2026; sources 60–69.** *Against the intake’s marked composition and the catalogue (checked at the write). Item II.6 records, as an unproved composition of sources 04 and 17, source 17’s retention $1/(2\eta L^m)$ under $|B| \geq 2(1 + \binom{m}{2})L^m$; this corollary proves it with the smaller threshold $2(1 + \binom{m}{2})L^{m-1}$, one power of L less, on every prime modulus. At $m = 16$, $2(1 + \binom{16}{2}) = 242$ and $(m+2)^m = 18^{16}$, as in Corollary II.330.6. The catalogue’s lemma_9_3 (Sections08_09.lean, line 77) quantifies over all moduli, so this prime-modulus corollary does not imply it; it is not formalized.*

The same conclusion holds for arbitrary finite abelian G, H whenever multiplication by $1, \dots, 2L-2$ is injective on G and every nonzero element of H has order greater than $L-1$. For example, composite cyclic groups whose prime factors exceed $2L-2$ satisfy both requirements when used as both domain and codomain. The general certificate in Corollary II.330.3 remains valid with small torsion present; its positivity is then an explicit additional condition.

II.330.4 Finite deterministic extraction

For a rational tolerance η , the existence proof can be made deterministic without an appeal to a limiting probability space. Enumerate the $|G||H|Q$ choices of (χ, ψ, j) and compute the conditional expectation of $\eta X - Y$ under the independent point selections. At least one choice has conditional expectation at least the unconditional expectation. Fix a maximizing choice.

Expose the membership of the points of B one at a time. At each step the current conditional expectation is a convex combination of the two conditional expectations obtained by including

or excluding the next point. Choose a branch with at least the current value. If its Bernoulli probability is 0 or 1, use the forced branch. Once every point is exposed, the resulting set has $\eta X - Y$ at least the proved lower bound.

All comparisons can be exact. The kernel probabilities belong to a cyclotomic field with conductor dividing

$$\text{lcm}(\exp(G), \exp(H), Q),$$

and their relevant real algebraic values can be compared exactly. Direct enumeration of additive tuples suffices to compute each conditional expectation; its survival factor is the product over *distinct* unexposed points. This supplies a finite extraction procedure. It is not a claim that the numerical thresholds or the straightforward implementation are practical at very small tolerances.

II.330.5 Sixteen-term relations and comparison of thresholds

The density-sensitive moment input used in sources 04 and 17 is valid on every finite abelian group. For completeness, if A is a nonempty finite subset of an abelian group and $E_j(A)$ counts its additive $2j$ -tuples, then

$$E_j(A) \geq \frac{E_2(A)^{j-1}}{|A|^{j-2}} \quad (j \geq 2). \quad (\text{II.330.20})$$

Using unnormalized Fourier coefficients on the ambient group, define a probability distribution on its dual with weights proportional to $|\widehat{1_A}(\chi)|^2$. Its $(j-1)$ st moment of $|\widehat{1_A}(\chi)|^2$ is $E_j(A)/|A|$, and its first moment is $E_2(A)/|A|$. Jensen's inequality proves (II.330.20).

Apply this to the graph $\{(x, \phi(x)) : x \in B\}$. If ϕ respects at least γb^3 additive quadruples, then it respects at least $\gamma^7 b^{15}$ additive sixteen-tuples. Consequently Corollary II.330.5 gives the following explicit form of the order-eight restriction step.

Corollary II.330.6 (Order eight). *Let $B \subseteq \mathbb{F}_p$ and $\phi : B \rightarrow \mathbb{F}_p$ respect at least γb^3 additive quadruples, where $0 < \gamma \leq 1$. For $0 < \eta \leq 1$, set*

$$L = \left\lceil \frac{17}{\gamma^7 \eta} \right\rceil.$$

If $b \geq 242L^{15}$, there is $B' \subseteq B$ for which at least a proportion $1/(1+\eta)$ of all additive sixteen-tuples is respected and

$$T_{16}(B') \geq \frac{b^{15}}{2\eta L^{16}} \geq \frac{\gamma^{112} \eta^{15}}{2 \cdot 18^{16}} b^{15}. \quad (\text{II.330.21})$$

The comparison uses the same normalization, the number of available points b , and the same L . Source 17's retained coefficient is $[2(m+2)^m]^{-1}$, exactly the coefficient in (II.330.18). The proposed combination of its Fejér signal with source 04's weighted rank-two estimate would require $2(1 + \binom{m}{2})L^m$ points. Our proof replaces that by $2(1 + \binom{m}{2})L^{m-1}$, a factor L smaller.

At $m = 16$ the two combined thresholds are therefore $242L^{16}$ and $242L^{15}$. Source 17's original exceptional-vector count has bandwidth degree 31 before its additional density dependence is removed, while source 04's order-eight theorem has threshold

$$242 \cdot 3200^8 (\gamma^7 \eta)^{-16}$$

and retention coefficient $(2 \cdot 3200^8)^{-1}$. The new theorem has degree 15 in the inverse parameter $\gamma^7 \eta$ and the larger retained coefficient $(2 \cdot 18^{16})^{-1}$; indeed $18^{16} = 324^8 < 3200^8$. These comparisons strengthen the stated local inputs of [244, 245], rather than an already completed global bound.

The threshold exponent $m - 1$ is not asserted optimal. On groups with persistent small-order torsion, a one-kernel selector can have substantial principal-line leakage: for fixed q , the ratio $S_{m,L}[q]/S_{m,L}$ tends to $1/q$ as $L \rightarrow \infty$. Removing such leakage, or obtaining a much lower threshold uniformly over growing prime cyclic groups, requires an additional argument. Formula (II.330.12) identifies precisely which relation statistics that argument must control.

II.331 Source 60's chapter: the arrangement polynomial of every order [write]

Added in the eleventh write, 7 October 2026.

The next six sections are source 60's Sections 3–8, on Gowers's respected arrangements of Section 12 near complete agreement on a full product. *Logically they continue source 33's chapter* (Section II.136), whose questions they answer in part; they are printed here because no number printed earlier may move. Source 60's introduction is Section F.65; its coarse decoder is its Appendix A, in Part V.

What is proved. For every even number $s \geq 4$ of vertical differences and distance $0 < \delta \leq 1/4$ to the exact models $a(x) + B(x, y) + \ell(y)$, the failure probability satisfies $\varepsilon_s \geq B_s(\delta) = (1 - (1 - 4\delta)^s)/4 - 2^{2s-3}\delta^{s-1}(1 - 2\delta)$, a polynomial that strictly improves source 33's for $s \geq 6$ and is source 33's cubic $4\delta - 24\delta^2 + 32\delta^3$ at $s = 4$ (Theorem II.334.1); equality for $0 < \delta < 1/4$ holds exactly for $a\mathbf{1}_C(x)\lambda(y)$ after adding an exact model, with $2a = 0 \neq a$, C a coset of index $n \geq 3$ and λ onto $\mathbb{Z}/2\mathbb{Z}$, at $\delta = 1/(2n)$ (Theorem II.334.3); under source 33's activation threshold the inverse branch gives every Taylor coefficient of the repair bound optimally in succession, $\delta \leq \varepsilon/16 + (15/128)\varepsilon^2 + (155/512)\varepsilon^3 + O(\varepsilon^4)$ for Gowers's eight-arrangements (Theorem II.335.2, Corollary II.335.3). The auxiliary-label profile behind it holds without conditions on the groups (Theorems II.333.2, II.333.3). Odd targets and the common prime: the bracket $(s-1)/s^2 \leq c_s^{\text{prime}} \leq (3s-2)/(2s^2)$, against source 33's $(3s-2)/(4s^2)$ and $2(s-1)/s^2$ (Theorems II.336.2, II.336.3). Two directions: if $\eta = \text{dist}(f, \mathcal{V}) + \text{dist}(f, \mathcal{H}) < 1/4$, then $\text{dist}(f, \mathcal{A}) \leq J(\eta) = (1 + \eta - \sqrt{1 - 2\eta})/2$, optimal (Theorem II.337.2), and biaffine repair with every inverse coefficient through degree $s - 2$ optimal (Theorems II.337.6, II.337.7). These answer source 33's Questions 2, 3 (the equality part) and 8 (locally), and improve both ends of its Question 1 (item II.49). The majority correction is Blum–Luby–Rubinfeld's and the coarse decoder source 33's, credited; the theorems concern full products, not the sparse domains Gowers needs (item II.25). Source 96, placed during this write and not yet written in, proves the odd-target repair and the bracket $[15/256, 23/256]$ independently, in a version with no cubic term and no smallness condition (from its placement record, not checked here; note after Theorem II.336.3).

Relation to the formal statements (checked at HEAD fb9602e55). `lemma_12_3` (Sections12_13.lean, line 151) has the exact companion `lemma_12_3_holds` (Proofs12HigherArrangements.lean, line 334); `lemma_12_5` and `lemma_12_6` (lines 170 and 183) have no counterpart of the repair theorem. Nothing in this chapter is formalized.

II.332 Respected arrangements and their normalization

Source 60, delivered Section 3.

Throughout this section, G, H, K are finite abelian groups and $s = 2m \geq 4$. All probabilities use

uniform measures, independently unless a relation is specified. For $f : G \times H \rightarrow K$, write

$$D_h f(x, y) = f(x, y + h) - f(x, y).$$

The vertical s -difference test chooses $h \in H$, independent base points $y_1, \dots, y_s \in H$, and a uniform tuple $(x_1, \dots, x_s)$ subject to

$$x_1 + \dots + x_m = x_{m+1} + \dots + x_s. \quad (\text{II.332.1})$$

It accepts precisely when

$$\sum_{i=1}^m D_h f(x_i, y_i) = \sum_{i=m+1}^s D_h f(x_i, y_i). \quad (\text{II.332.2})$$

Its failure probability is $\varepsilon_s(f)$. The ordered-query convention includes $h = 0$, repeated coordinates, and coinciding vertices. There are exactly

$$|G|^{s-1} |H|^{s+1} \quad (\text{II.332.3})$$

query tuples. If $G = H = \mathbb{Z}/N\mathbb{Z}$, this is N^{2s} . In the arrangement terminology of Gowers's Section 12, $s = 2k$ corresponds to a k -arrangement; in particular, eight-arrangements have $s = 16$. These are the conventions of the earlier full-product repair report [248], derived from Gowers's framework [1].

Let

$$\mathcal{V}(G, H; K) = \{(x, y) \mapsto a(x) + B(x, y) + \ell(y) : B \text{ biadditive, } \ell \text{ additive}\}, \quad (\text{II.332.4})$$

where $a : G \rightarrow K$ is arbitrary. Uniform Hamming distance is denoted by d , and

$$\delta = \delta_{\mathcal{V}}(f) := \min_{\psi \in \mathcal{V}} d(f, \psi).$$

Every member of $\mathcal{V}$ passes every test: its derivative is $B(x, h) + \ell(h)$, and (II.332.1) cancels the biadditive terms. Thus adding a member of $\mathcal{V}$ changes neither the test outcome nor the distance to $\mathcal{V}$. The converse, that a function passing every test belongs to $\mathcal{V}$, follows from the coarse decoder reproduced in Appendix V.227.

Lemma II.332.1 (Independence of proper coordinate families). *Every proper subfamily of the coordinates in a uniform tuple satisfying (II.332.1) is independent and uniform. The same holds for the pairs (x_i, y_i) , conditional on any fixed increment h .*

Proof. After prescribing any proper coordinate family, choose one unprescribed coordinate to solve for. All other unprescribed coordinates are free, and the coefficient of the final coordinate is 1 or -1 . There are therefore exactly $|G|^{s-|I|-1}$ completions after prescribing the coordinates in I , independently of their values. The labels y_i are independently uniform by definition. $\square$

The preceding independence is almost complete, but the omitted last degree matters for sharp bounds. The next theorem retains it through Fourier coefficients instead of discarding arrangements with several erroneous differences.

II.333 A sharp spectral inequality with independent labels

Source 60, delivered Section 4.

Let S be a finite nonempty set and $F : G \times S \rightarrow K$. Choose the horizontal coordinates as in (II.332.1) and labels $u_1, \dots, u_s \in S$ independently. Write $e_s(F)$ for the probability that

$$\sum_{i \leq m} F(x_i, u_i) \neq \sum_{i > m} F(x_i, u_i).$$

This auxiliary test does not itself involve derivatives.

For a character $\chi \in \widehat{K}$, put

$$g_\chi(x) = \mathbb{E}_{u \in S} \chi(F(x, u)), \quad \widehat{g}_\chi(\xi) = \mathbb{E}_{x \in G} g_\chi(x) \overline{\xi(x)}.$$

The horizontal Fourier transform is normalized; the sum over $\widehat{G}$ below is not normalized. The average over $\widehat{K}$ is normalized.

Lemma II.333.1 (Moment identity). *For every even $s \geq 4$,*

$$1 - e_s(F) = \mathbb{E}_{\chi \in \widehat{K}} \sum_{\xi \in \widehat{G}} |\widehat{g}_\chi(\xi)|^s. \quad (\text{II.333.1})$$

Proof. Character orthogonality on K expresses acceptance as the character average of the signed sum. Averaging the independent labels leaves $g_\chi(x_i)$ on the first half and its complex conjugate on the second. Expand each function in horizontal characters. Orthogonality in the $s - 1$ free coordinates forces one common frequency in all s factors, giving (II.333.1). $\square$

Define, for $0 \leq t \leq 1/2$,

$$E_s(t) = \frac{1 - (1 - 2t)^s}{2} - 2^{s-1} t^{s-1} (1 - t). \quad (\text{II.333.2})$$

Theorem II.333.2 (Auxiliary-label profile). *If $t = \mathbb{P}_{x,u}[F(x, u) \neq 0] \leq 1/2$, then*

$$e_s(F) \geq E_s(t). \quad (\text{II.333.3})$$

No condition on the exponent of any group or on the number of labels is required.

Proof. The case $t = 0$ is immediate. Suppose $t > 0$, put $q = 2t$ and $b = (1 - q)^2$, and write

$$a_\chi = \mathbb{E}_{x,u} \chi(F(x, u)), \quad v_\chi = |a_\chi|^2.$$

Because $F = 0$ with probability $1 - t$,

$$|a_\chi - 1| \leq 2t, \quad b \leq v_\chi \leq 1. \quad (\text{II.333.4})$$

For a nontrivial horizontal character,

$$\widehat{g}_\chi(\xi) = \mathbb{E}_{x,u} \mathbf{1}_{\{F \neq 0\}} (\chi(F(x, u)) - 1) \overline{\xi(x)},$$

so $|\widehat{g}_\chi(\xi)| \leq q$. Parseval and $|g_\chi| \leq 1$ give

$$\begin{aligned} \sum_{\xi \neq 1} |\widehat{g}_\chi(\xi)|^s &\leq q^{s-2} \sum_{\xi \neq 1} |\widehat{g}_\chi(\xi)|^2 \\ &\leq q^{s-2}(1 - v_\chi). \end{aligned} \quad (\text{II.333.5})$$

Set

$$C = \frac{1 - b^m}{1 - b} = \sum_{j=0}^{m-1} b^j.$$

The chord of the convex function $z \mapsto z^m$ on $[b, 1]$ gives

$$v_\chi^m \leq 1 - C(1 - v_\chi). \quad (\text{II.333.6})$$

Combining the last two inequalities yields

$$1 - \sum_{\xi} |\widehat{g}_\chi(\xi)|^s \geq (C - q^{s-2})(1 - v_\chi).$$

The factor $C - q^{s-2}$ is nonnegative for $0 < q \leq 1$, and is positive when $q < 1$. If $P(z) = \mathbb{P}[F = z]$, character orthogonality gives

$$\mathbb{E}_\chi v_\chi = \sum_z P(z)^2 \leq (1 - t)^2 + t^2. \quad (\text{II.333.7})$$

Hence the moment identity implies

$$\begin{aligned} e_s(F) &\geq 2t(1 - t)(C - q^{s-2}) \\ &= \frac{1 - (1 - 2t)^s}{2} - 2^{s-1}t^{s-1}(1 - t), \end{aligned}$$

as required. At $t = 1/2$, the same calculation has $b = 0$, $C = 1$, and gives the valid endpoint bound zero. $\square$

Theorem II.333.3 (All interior equality cases). *Suppose $0 < t < 1/2$. Equality in (II.333.3) holds if and only if*

$$F(x, u) = a\mathbf{1}_C(x), \quad (\text{II.333.8})$$

where $a \in K$ is a nonzero element of order two and C is a coset of a subgroup of G , with density t . In particular, the labels have no effect and t is the reciprocal of an integer at least three.

Proof. All slack terms in the preceding proof are nonnegative. Since $C - q^{s-2} > 0$, equality forces equality in (II.333.7). The mass outside zero must therefore be concentrated at one nonzero element a . Thus F takes only the values $0, a$.

The chord inequality (II.333.6) is strict in the interior of $[b, 1]$, because $m \geq 2$. Consequently

$$|1 - t + t\chi(a)|^2 \in \{(1 - 2t)^2, 1\} \quad (\chi \in \widehat{K}).$$

This is equivalent to $\chi(a) \in \{1, -1\}$ for every character. Characters separate points, so $2a = 0$; since $a \neq 0$, its order is two. Choose a character with $\chi(a) = -1$.

For this character, equality in the final step of (II.333.5) forces $|g_\chi(x)| = 1$ for every x : the multiplier q^{s-2} is positive. But $g_\chi(x)$ is the average of the signs distinguishing $F(x, u) = 0$ and

$F(x, u) = a$. Thus these signs, and hence $F(x, u)$, are independent of u at each x . Let C be the nonzero support.

Equality in the first step of (II.333.5) now says that every nontrivial Fourier coefficient of $\mathbf{1}_C$ has modulus either zero or t . The trivial coefficient has modulus t as well. Fix $c \in C$ and let J be the subgroup generated by $C - C$. Equality in the triangle inequality for $\widehat{\mathbf{1}_C}(\xi)$ says that its modulus is t exactly when ξ is constant on C , equivalently when $\xi \in J^\perp$. Therefore its Fourier support is precisely $J^\perp$. Fourier inversion gives

$$\mathbf{1}_C(x) = t \sum_{\xi \in J^\perp} \xi(x - c) = t|J^\perp| \mathbf{1}_{c+J}(x).$$

Evaluation at c proves $t|J^\perp| = 1$ and $C = c + J$.

Conversely, for a function of the form (II.333.8), half the characters of K satisfy $\chi(a) = 1$ and give $g_\chi = 1$. The other half give $g_\chi = 1 - 2\mathbf{1}_C$. If $n = [G : J] = 1/t$, this function has constant Fourier coefficient $1 - 2t$ and $n - 1$ other nonzero coefficients, each of modulus $2t$. Thus

$$e_s(F) = \frac{1}{2} \{1 - (1 - 2t)^s - (n - 1)(2t)^s\} = E_s(t),$$

which proves sufficiency. $\square$

Lemma II.333.4 (Monotone profile ratio). *The ratio $E_s(t)/t$, extended at zero by the value s , is strictly decreasing on $[0, 1/2]$. Its value at $1/2$ is zero. In particular, $E_s(t) > 0$ for $0 < t < 1/2$.*

Proof. For positive t ,

$$\frac{E_s(t)}{t} = \sum_{j=0}^{s-1} (1 - 2t)^j - 2^{s-1} t^{s-2} (1 - t).$$

The first term is strictly decreasing; its $j = 1$ summand already is. The derivative of $t^{s-2}(1 - t)$ is $t^{s-3}((s - 2) - (s - 1)t) > 0$ on $(0, 1/2]$. This proves strict decrease. The endpoint values follow by substitution and continuity. $\square$

II.334 The complete polynomial and its equality cases

Source 60, delivered Section 5.

Choose a nearest model ψ and put $r = f - \psi$. Its arbitrary row constants can be optimized separately. We may therefore suppose zero is a most frequent value of $r(x, \cdot)$ in every row. Define

$$\begin{aligned} p_x(z) &= \mathbb{P}_y[r(x, y) = z], & \delta_x &= 1 - p_x(0), \\ t_h &= \mathbb{P}_{x,y}[D_h r(x, y) \neq 0], & \bar{t} &= \mathbb{E}_h t_h. \end{aligned} \tag{II.334.1}$$

Then

$$\delta = \mathbb{E}_x \delta_x, \quad 0 \leq t_h \leq 2\delta, \quad \bar{t} = \mathbb{E}_x \left(1 - \sum_z p_x(z)^2 \right) \geq \delta. \tag{II.334.2}$$

Indeed, a nonzero derivative has at least one nonzero endpoint. For the average, $(y, y + h)$ is uniformly independent, and $\sum_z p_x(z)^2 \leq \max_z p_x(z) = p_x(0)$. If $\delta_x \leq \kappa \leq 1/2$ for every row, one also has

$$\bar{t} \geq 2(1 - \kappa)\delta, \tag{II.334.3}$$

because $\sum_z p_x(z)^2 \leq (1 - \delta_x)^2 + \delta_x^2$.

Define

$$\begin{aligned} B_s(d) &= \frac{1}{2} E_s(2d) \\ &= \frac{1 - (1 - 4d)^s}{4} - 2^{2s-3} d^{s-1} (1 - 2d) \\ &= \frac{1 - (1 - q)^s - 2q^{s-1} + q^s}{4}, \quad q = 4d. \end{aligned} \tag{II.334.4}$$

The terms of degree s cancel, so B_s has degree $s - 1$.

Theorem II.334.1 (Sharp all-order local polynomial). *For every even $s \geq 4$ and every f with $0 < \delta \leq 1/4$,*

$$\varepsilon_s(f) \geq \frac{\bar{t}}{\delta} B_s(\delta) \geq B_s(\delta). \tag{II.334.5}$$

Under the row cap $\delta_x \leq \kappa \leq 1/2$, the stronger conclusion is

$$\varepsilon_s(f) \geq 2(1 - \kappa) B_s(\delta). \tag{II.334.6}$$

The unprofiled bound also holds at $\delta = 0$.

Proof. At each fixed increment, apply Theorem II.333.2 to $F_h(x, y) = D_{hr}(x, y)$ with label set H . Since $t_h \leq 2\delta \leq 1/2$, Lemma II.333.4 implies

$$\varepsilon_s(f) \geq \mathbb{E}_h E_s(t_h) \geq \frac{E_s(2\delta)}{2\delta} \bar{t} = \frac{\bar{t}}{\delta} B_s(\delta).$$

Now use (II.334.2) or (II.334.3). □

For orientation, the first two polynomials are

$$B_4(d) = 4d - 24d^2 + 32d^3, \tag{II.334.7}$$

$$B_6(d) = 6d - 60d^2 + 320d^3 - 960d^4 + 1024d^5. \tag{II.334.8}$$

The first is the sharp cubic already proved in [248]. The second and the higher polynomials close the gap left by that report's one-error estimate. More explicitly, its conditional polynomial is

$$L_s(t) = st(1 - t)((1 - t)^{s-2} - t^{s-2}),$$

and direct parity expansion gives

$$E_s(t) - L_s(t) = \sum_{\substack{3 \leq j \leq s-3 \\ j \text{ odd}}} \binom{s}{j} \{t^j(1 - t)^{s-j} - t^{s-1}(1 - t)\}. \tag{II.334.9}$$

Every summand is positive for $0 < t < 1/2$. The sum is empty for $s = 4$ and nonempty for every even $s \geq 6$. Therefore the new polynomial strictly improves the earlier one throughout $0 < d < 1/4$ in those higher orders.

Lemma II.334.2 (Distance of a single coset spike to affine maps). *Let $J \leq G$ have index $n \geq 3$, let C be a coset of J , and let v be a nonzero element of an abelian group M . The distance of $x \mapsto v\mathbf{1}_C(x)$ to the affine maps $G \rightarrow M$ is exactly $1/n$.*

Proof. The zero map gives the upper bound. Consider an affine map Λ with additive part L . If $L|_J$ is nonzero, the proportion matching any constant target in each J -coset is at most $1/2$, because a nonempty fiber of $L|_J$ has index at least two. Thus the disagreement is at least $1/2$.

If $L|_J = 0$, descend to G/J , whose order is n . A constant affine map cannot improve on zero. For a nonconstant affine map, let $r \geq 2$ be the size of its additive image; then r divides n . It matches zero on at most n/r quotient points and the exceptional nonzero target on at most one further point. If $n \geq 4$, $n/r + 1 \leq n/2 + 1 \leq n - 1$. If $n = 3$, necessarily $r = 3$, and the same bound is $n/r + 1 = 2 = n - 1$. At least one quotient point is therefore wrong, giving disagreement at least $1/n$. $\square$

Theorem II.334.3 (Classification of equality in the local polynomial). *Let $s \geq 4$ be even and $0 < \delta_{\mathcal{V}}(f) < 1/4$. Then*

$$\varepsilon_s(f) = B_s(\delta_{\mathcal{V}}(f)) \quad (\text{II.334.10})$$

if and only if there exist $\psi \in \mathcal{V}$, a nonzero element $a \in K$ of order two, a surjective homomorphism $\lambda : H \rightarrow \mathbb{Z}/2\mathbb{Z}$, and a coset C of a subgroup $J \leq G$ with index $n \geq 3$, such that

$$f(x, y) = \psi(x, y) + a\mathbf{1}_C(x)\lambda(y). \quad (\text{II.334.11})$$

The last product denotes 0 or a according to the binary values. Every such function has

$$\delta_{\mathcal{V}}(f) = \frac{1}{2n}, \quad (\text{II.334.12})$$

and attains equality simultaneously for every even $s \geq 4$.

Proof. First assume equality and choose the modal nearest residual r used above. Since $B_s(\delta) > 0$ and the ratio in Lemma II.333.4 is strictly decreasing, every step in (II.334.5) can be an equality only if

$$\bar{t} = \delta, \quad t_h \in \{0, 2\delta\} \quad \text{for every } h. \quad (\text{II.334.13})$$

Moreover, equality holds in Theorem II.333.2 at each active increment, meaning each h with $t_h > 0$. By Theorem II.333.3, $D_h r(x, y)$ is independent of y at every active increment. It is zero everywhere at every inactive increment. Thus every derivative of each row is independent of its base point, and

$$r(x, y) = c(x) + L_x(y)$$

for homomorphisms $L_x : H \rightarrow K$.

The common kernel

$$H_0 = \{h : L_x(h) = 0 \text{ for every } x\}$$

is exactly the set of inactive increments. Equation (II.334.13) says that the active proportion is $1/2$, so $[H : H_0] = 2$. Let $\lambda : H \rightarrow \mathbb{Z}/2\mathbb{Z}$ have kernel H_0 . Every L_x factors through λ . Fix one active increment h_0 . Theorem II.333.3 gives

$$L_x(h_0) = a\mathbf{1}_C(x),$$

where a has order two and C is a subgroup coset of density 2δ . Consequently $L_x(y) = a\mathbf{1}_C(x)\lambda(y)$. Absorbing $c(x)$ into the model proves (II.334.11). The index is greater than two because $\delta < 1/4$.

For the converse, it suffices to treat $f_0(x, y) = a\mathbf{1}_C(x)\lambda(y)$, since adding a model changes neither distance nor test outcomes. Compare f_0 with an arbitrary model $b(x) + B(x, y) + \ell(y)$. Its row slope is the affine map

$$\Lambda : G \longrightarrow \text{Hom}(H, K), \quad \Lambda(x) = B(x, \cdot) + \ell.$$

Lemma II.334.2, applied to the nonzero homomorphism $a\lambda$, says that these slopes disagree with those of f_0 on at least $1/n$ of the rows. In every such row, the difference is a nonconstant affine map of H , whose zero set has density at most $1/2$. Even after optimizing all row constants, the distance is therefore at least $1/(2n)$. The zero model attains that distance.

For half the increments, $\lambda(h) = 0$ and the test always accepts. For the other half, $D_h f_0(x, y) = a\mathbf{1}_C(x)$, and Theorem II.333.3 gives failure $E_s(1/n)$. Averaging proves $\varepsilon_s(f_0) = E_s(1/n)/2 = B_s(1/(2n))$. $\square$

This classification includes even subgroup indices. The earlier example in [248] used odd cyclic orders to simplify the nearest-model argument. Lemma II.334.2 removes that restriction. It also shows the exact scope of the word sharp: the polynomial is attained at the classified distances $1/(2n)$, including a sequence tending to zero. The theorem does not claim that it is the best pointwise bound at every other feasible distance.




Figure II.334.1: Left: the new $s = 16$ polynomial and the earlier F_{16} on $0 \leq \delta \leq 1/4$. Markers are equality examples at $\delta = 1/(2n)$, $3 \leq n \leq 24$; the shaded interval is the increasing branch used below. The full polynomial is not globally increasing. Right: a classified equality map on $\mathbb{Z}_{12} \times \mathbb{Z}_8$ with target $\mathbb{Z}_2$, $C = \{0, 4, 8\}$ and $\ell_2(y) = y \bmod 2$. The dark cells have value one. Its distance is $1/8$, illustrating an even subgroup index.

II.335 A repair modulus sharp to every asymptotic order

Source 60, delivered Section 6.

To invert a local polynomial one must first locate the nearest model on its small branch. We use the previously established elementary decoder from [248], whose proof is included in Appendix V.227:

$$\varepsilon_4(f) < \frac{1}{216} \implies \delta_V(f) \leq 60\varepsilon_4(f). \quad (\text{II.335.1})$$

The moment identity also implies $\varepsilon_4(f) \leq \varepsilon_s(f)$: every Fourier coefficient has modulus at most one, so the acceptance moment decreases with s .

Lemma II.335.1 (An explicit increasing branch). *The polynomial B_s is strictly increasing on $[0, 1/(4(s-1))]$, and*

$$B_s\left(\frac{1}{4(s-1)}\right) > \frac{23}{216}. \quad (\text{II.335.2})$$

Proof. With $q = 4d$,

$$B'_s(d) = s(1-q)^{s-1} - 2(s-1)q^{s-2} + sq^{s-1}.$$

For $0 \leq q \leq 1/(s-1)$ this is at least

$$s\left(\frac{s-2}{s-1}\right)^{s-1} - \frac{2}{(s-1)^{s-3}} > 0.$$

The final inequality is equivalent to $s(s-2)^{s-1} > 2(s-1)^2$. For $s \geq 4$ it follows from $(s-2)^{s-1} \geq (s-2)^3$, $s(s-2) > 2(s-1)$, and $(s-2)^2 \geq s-1$.

At the endpoint $q = 1/(s-1)$, Bernoulli's inequality gives $(1-q)^s \leq 1/(1+sq) < 1/2$. Also $2q^{s-1} \leq 2/27$. Substitution in (II.334.4), retaining the nonnegative q^s term or discarding it, proves

$$B_s\left(\frac{1}{4(s-1)}\right) > \frac{1}{4}\left(\frac{1}{2} - \frac{2}{27}\right) = \frac{23}{216}.$$

□

Let ρ_s denote the inverse branch of B_s with values in $[0, 1/(4(s-1))]$.

Theorem II.335.2 (All-order repair). *If*

$$\varepsilon_s(f) \leq \frac{1}{240(s-1)}, \quad (\text{II.335.3})$$

then

$$\delta_V(f) \leq \rho_s(\varepsilon_s(f)). \quad (\text{II.335.4})$$

The activation threshold is the same sufficient threshold as in [248]; no optimality of that threshold is asserted.

Proof. The moment comparison and (II.335.1) give

$$\delta \leq 60\varepsilon_s(f) \leq \frac{1}{4(s-1)}.$$

The coarse decoder applies because $1/(240(s-1)) \leq 1/720 < 1/216$. Theorem II.334.1 gives $B_s(\delta) \leq \varepsilon_s(f)$. Lemma II.335.1 places this error in the domain of ρ_s and permits inversion on the branch containing δ . □

Corollary II.335.3 (Optimal higher coefficients). *For every even $s \geq 6$,*

$$\rho_s(u) = \frac{u}{s} + \frac{2(s-1)}{s^2}u^2 + \frac{8(s-1)(2s-1)}{3s^3}u^3 + O_s(u^4). \quad (\text{II.335.5})$$

For $s = 4$,

$$\rho_4(u) = \frac{u}{4} + \frac{3}{8}u^2 + u^3 + \frac{105}{32}u^4 + O(u^5). \quad (\text{II.335.6})$$

Every Taylor coefficient of ρ_s is optimal in succession for uniform asymptotic repair estimates over finite abelian groups.

Proof. For $s \geq 6$, the correction term in (II.334.4) begins at degree $s - 1 \geq 5$, so

$$B_s(d) = sd - 2s(s-1)d^2 + \frac{8}{3}s(s-1)(s-2)d^3 + O_s(d^4).$$

Substituting a power series into $B_s(\rho_s(u)) = u$ yields (II.335.5). The case $s = 4$ instead uses the whole cubic (II.334.7), giving (II.335.6).

The inverse is analytic near zero because $B'_s(0) = s > 0$. For each $n \geq 3$, take $G = \mathbb{Z}/n\mathbb{Z}$, $H = K = \mathbb{Z}/2\mathbb{Z}$, $C = \{0\}$, and $f_n(x, y) = \mathbf{1}_C(x)y$. Theorem II.334.3 gives

$$\delta_n = \frac{1}{2n}, \quad \varepsilon_s(f_n) = B_s(\delta_n).$$

For sufficiently large n , these examples lie on the inverse branch and satisfy $\delta_n = \rho_s(\varepsilon_s(f_n))$ exactly. Since their errors tend to zero, reducing any coefficient while keeping all earlier coefficients fixed contradicts these equalities at the first changed order. The statement allows the remainder constant to depend on fixed s , but not on the groups or the function. $\square$

For example,

$$\begin{aligned} \rho_6(u) &= \frac{u}{6} + \frac{5}{18}u^2 + \frac{55}{81}u^3 + O(u^4), \\ \rho_{16}(u) &= \frac{u}{16} + \frac{15}{128}u^2 + \frac{155}{512}u^3 + O(u^4). \end{aligned}$$

At $s = 16$, the quadratic inverse in [248] has cubic coefficient $225/512$; the sharp value is $155/512$. More generally, classical Lagrange inversion gives a reproducible formula for every coefficient:

$$[u^j]\rho_s(u) = \frac{1}{j}[z^{j-1}] \left(\frac{z}{B_s(z)} \right)^j. \quad (\text{II.335.7})$$

Thus the higher-arrangement coefficient question in [248] is answered to every order. The exact equality portion of its cubic classification question is answered by Theorem II.334.3; quantitative near-equality is a separate problem.

II.336 Odd targets and a sharper common-prime coefficient interval

Source 60, delivered Section 7.

The equality family uses two-torsion, so it does not determine the optimal second coefficient when $G = H = K = \mathbb{F}_p$ with p odd. The earlier report [248] placed that optimal coefficient between $(3s-2)/(4s^2)$ and $2(s-1)/s^2$. We improve both ends. An expansion through two erroneous queries gives a stronger repair upper bound whenever the target has odd order. A block of rows carrying a two-valued interval function gives a larger lower obstruction over a common prime field. The resulting interval is still not a determination of the optimal coefficient.

Lemma II.336.1 (A uniform expansion through two errors). *Let $Z_1, \dots, Z_s$ be K -valued random variables, where $s = 2m \geq 4$. Suppose every three-coordinate subfamily is independent, each marginal is zero with probability $1 - t$, and the conditional distribution outside zero is q . Put*

$$c(q) = \sum_{z \neq 0} q(z)^2, \quad d(q) = \sum_{z \neq 0} q(z)q(-z), \quad D_s = (s^2 + 1) \binom{s}{3}.$$

For m positive signs and m negative signs, the failure probability $e = \mathbb{P}[\sum_i \sigma_i Z_i \neq 0]$ satisfies

$$\left| e - st + \left\{ \binom{s}{2} + m^2 c(q) + m(m-1)d(q) \right\} t^2 \right| \leq D_s t^3. \quad (\text{II.336.1})$$

If $t = 0$, the conclusion is interpreted as $e = 0$.

Proof. For $t > 0$, put $I_i = \mathbf{1}_{\{Z_i \neq 0\}}$ and $C_{ij} = \mathbf{1}_{\{\sigma_i Z_i + \sigma_j Z_j = 0\}}$. The expression

$$T = \sum_i I_i - \sum_{i < j} (1 + C_{ij}) I_i I_j$$

equals the failure indicator whenever at most two variables are nonzero. Opposite-sign pairs cancel with conditional probability $c(q)$; there are m^2 such pairs. Same-sign pairs cancel with conditional probability $d(q)$; there are $m(m-1)$ of them. Pairwise independence therefore gives

$$\mathbb{E}T = st - \left\{ \binom{s}{2} + m^2 c(q) + m(m-1)d(q) \right\} t^2.$$

At every outcome, the absolute discrepancy between T and the failure indicator is at most $s^2 + 1$. The probability of at least three nonzero variables is at most $\binom{s}{3} t^3$, by the union bound and three-coordinate independence. This proves the claim. $\square$

Theorem II.336.2 (A stronger quadratic repair coefficient for odd targets). *Suppose K has odd order, and set*

$$C_s = \frac{s(3s-2)}{4}, \quad Q_s(d) = sd - 2C_s d^2 - 4D_s d^3, \quad D_s = (s^2 + 1) \binom{s}{3}.$$

For the auxiliary test of Section II.333, without a smallness assumption on $t = \mathbb{P}[F \neq 0]$,

$$e_s(F) \geq st - C_s t^2 - D_s t^3. \quad (\text{II.336.2})$$

For arrangements, if $\delta = \delta_V(f) \leq 1/(4s^2)$, then $\varepsilon_s(f) \geq Q_s(\delta)$. If

$$\varepsilon_s(f) \leq \frac{1}{240s^2}, \quad (\text{II.336.3})$$

the inverse branch of Q_s on $[0, 1/(4s^2)]$ gives the uniform bound

$$\delta_V(f) \leq \frac{\varepsilon_s(f)}{s} + \frac{3s-2}{2s^2} \varepsilon_s(f)^2 + O_s(\varepsilon_s(f)^3). \quad (\text{II.336.4})$$

No claim of sharpness is made for the quadratic coefficient in (II.336.4) or for the convenient cubic remainder and activation threshold.

Proof. Because K has no nonzero element of order two, its nonzero elements partition into distinct pairs $\{z, -z\}$. Therefore

$$c(q) + d(q) = \sum_{\{z, -z\}} (q(z) + q(-z))^2 \leq 1,$$

and

$$m^2 c(q) + m(m-1)d(q) \leq m^2.$$

Lemma II.336.1, with $\binom{s}{2} + m^2 = C_s$, proves (II.336.2).

For arrangements, take the modal nearest residual and the profile (II.334.1). Apply (II.336.2) at each increment. Since $t_h \leq 2\delta$,

$$\begin{aligned}\varepsilon_s(f) &\geq s\bar{t} - C_s\mathbb{E}_h t_h^2 - D_s\mathbb{E}_h t_h^3 \\ &\geq (s - 2C_s\delta - 4D_s\delta^2)\bar{t}.\end{aligned}$$

The multiplier is positive in the stated range, so $\bar{t} \geq \delta$ gives $\varepsilon_s(f) \geq Q_s(\delta)$. For completeness, both this positivity and the increasing inverse branch follow from one estimate. Since $C_s \leq 3s^2/4$ and $D_s \leq s^5/3$, for $0 \leq d \leq 1/(4s^2)$ one has

$$\begin{aligned}Q'_s(d) &= s - 4C_s d - 12D_s d^2 \\ &\geq s - \frac{3}{4} - \frac{s}{4} \geq \frac{9s}{16} > \frac{s}{2}.\end{aligned}$$

The multiplier above is at least $Q'_s(\delta)$.

Under (II.336.3), the coarse decoder gives $\delta \leq 60\varepsilon_s(f) \leq 1/(4s^2)$. Its own threshold is automatically satisfied. Moreover, $Q_s(1/(4s^2)) \geq 1/(8s) > 1/(240s^2)$, so the inverse is defined on the whole required error interval. Inverting Q_s gives the first two coefficients $1/s$ and $2C_s/s^3 = (3s-2)/(2s^2)$, proving (II.336.4) with a uniform remainder. $\square$

Theorem II.336.3 (Improved common-prime obstruction). *Fix an even $s \geq 4$. There are functions $f_p : \mathbb{F}_p^2 \rightarrow \mathbb{F}_p$, with p tending to infinity through odd primes and $\delta_p = \delta_V(f_p) \rightarrow 0$, such that*

$$\varepsilon_s(f_p) = s\delta_p - s(s-1)\delta_p^2 + O_s(\delta_p^3). \quad (\text{II.336.5})$$

Consequently, every uniform common-prime repair estimate of the form

$$\delta_V(f) \leq \frac{\varepsilon_s(f)}{s} + c\varepsilon_s(f)^2 + O_s(\varepsilon_s(f)^3),$$

with a remainder constant independent of p and f , must have $c \geq (s-1)/s^2$. If c_s^{prime} denotes the infimum of admissible coefficients in such uniform estimates, then

$$\frac{s-1}{s^2} \leq c_s^{\text{prime}} \leq \frac{3s-2}{2s^2}. \quad (\text{II.336.6})$$

The upper bound in this bracket is the odd-target theorem; the lower bound is the new obstruction. Both improve the corresponding endpoints recorded in [248].

Proof. Let $R = \lfloor \sqrt{p} \rfloor$, choose any $S_p \subset \mathbb{F}_p$ of size R , and put

$$A_p = \{0, 1, \dots, (p-3)/2\}, \quad f_p(x, y) = \mathbf{1}_{S_p}(x)\mathbf{1}_{A_p}(y).$$

The indicators are regarded as elements of $\mathbb{F}_p$. Write $r = R/p$ and $\alpha = (p-1)/(2p)$.

For zero slope, each row is best approximated by the constant zero: in the active rows, $\alpha < 1/2$. Its distance is $r\alpha$. To rule out other slopes, note that every model over the common prime field has the form $a(x) + (bx+c)y$. If $(b, c) \neq (0, 0)$, its slope is nonzero in at least $p-1$ rows, and on each such row its distance from zero is at least $1 - 1/p$, independently of $a(x)$. Thus that model's

distance from zero is at least $(1 - 1/p)^2$. By the triangle inequality, its distance from f_p is at least $(1 - 1/p)^2 - r\alpha$, which exceeds $r\alpha$ for all sufficiently large p . Therefore

$$\delta_p = r\alpha = \frac{R(p-1)}{2p^2}. \quad (\text{II.336.7})$$

Use representatives $h = 0, \dots, p-1$ and define

$$\rho_h = \mathbb{P}_y[\mathbf{1}_{A_p}(y+h) \neq \mathbf{1}_{A_p}(y)] = \frac{2 \min(h, p-h)}{p}.$$

The two nonzero derivative values $1, -1$ occur equally often. For fixed h , a uniformly sampled derivative of f_p therefore has probabilities $1 - r\rho_h, r\rho_h/2, r\rho_h/2$ at $0, 1, -1$. Lemma II.332.1 gives three-coordinate independence, even when $s = 4$. Here the nonzero conditional law assigns probability $1/2$ to each of $1, -1$, so $c(q) = d(q) = 1/2$. Applying Lemma II.336.1 conditionally and then averaging yields

$$\left| \varepsilon_s(f_p) - sr\mathbb{E}_h\rho_h + \frac{3}{4}s(s-1)r^2\mathbb{E}_h\rho_h^2 \right| \leq (s^2+1) \binom{s}{3} r^3. \quad (\text{II.336.8})$$

The error constant here is independent of p and of the choice of S_p .

The interval has two shifts of each length $1, \dots, (p-1)/2$, so the elementary sums of the first powers and squares give the exact moments

$$\mathbb{E}_h\rho_h = \frac{p^2-1}{2p^2}, \quad \mathbb{E}_h\rho_h^2 = \frac{p^2-1}{3p^2}. \quad (\text{II.336.9})$$

Substitution, followed by (II.336.7), gives

$$\varepsilon_s(f_p) = s\delta_p \left(1 + \frac{1}{p}\right) - s(s-1)\delta_p^2 \frac{p+1}{p-1} + O_s(r^3). \quad (\text{II.336.10})$$

Since $R = \lfloor \sqrt{p} \rfloor$, one has $r = O(\delta_p)$ and $p^{-1} = O(\delta_p^2)$. Thus (II.336.10) implies (II.336.5), with a uniform remainder.

Reversing the expansion proves

$$\delta_p = \frac{\varepsilon_s(f_p)}{s} + \frac{s-1}{s^2} \varepsilon_s(f_p)^2 + O_s(\varepsilon_s(f_p)^3),$$

which forces the lower bound in (II.336.6). Theorem II.336.2 supplies its upper bound. □

[write] 7 October 2026; sources 60–69. *A later independent proof* (forward pointer, not checked at this write). Source 96, placed by 09f5b745a during this write and not yet written in, proves the odd-target repair and the common-prime bracket $[15/256, 23/256]$ at $s = 16$ independently, in a version with no cubic term and no smallness condition, according to its placement record; it is to be written with source 96. This source arrived first, and the results are printed once, in its form.

The growing number of active rows is essential to the second-order calculation. With a single active row, the unavoidable odd-prime imbalance between the two row values contributes at second order. Here it contributes only to the uniform cubic remainder. For eight-arrangements the improved bracket is

$$\frac{15}{256} \leq c_{16}^{\text{prime}} \leq \frac{23}{256},$$

in place of the earlier interval $[23/512, 15/128]$. Determining the optimal value inside the new interval remains open.

II.337 Two-direction testing and sharp robust intersection

Source 60, delivered Section 8.

The vertical test is invariant under adding an arbitrary function of x . Consequently a vertical repair theorem cannot, by itself, produce a fully biaffine map. The horizontal analogue removes this freedom, but only after the two approximate models have been synchronized. The following argument gives an exact robust-intersection modulus before bringing in the arrangement tests.

II.337.1 The model classes and their separation

For finite abelian G, H, K , let

$$\begin{aligned}\mathcal{V} &= \{a(x) + B(x, y) + \ell(y)\}, \\ \mathcal{H} &= \{b(y) + C(x, y) + m(x)\}, \\ \mathcal{A} &= \{c + D(x, y) + m(x) + \ell(y)\},\end{aligned}$$

where $a : G \rightarrow K$ and $b : H \rightarrow K$ are arbitrary, B, C, D are biadditive, $m : G \rightarrow K$ and $\ell : H \rightarrow K$ are homomorphisms, and $c \in K$. Clearly $\mathcal{A} \subseteq \mathcal{V} \cap \mathcal{H}$. The proof below also establishes the reverse inclusion.

Lemma II.337.1 (Separation of different bilinear parts). *Suppose $v \in \mathcal{V}$ and $w \in \mathcal{H}$ have respective biadditive parts B and C . If $B \neq C$, then*

$$\text{dist}(v, w) \geq \frac{1}{4}.$$

The constant $1/4$ cannot be increased in general.

Proof. Write $D = B - C$ and $v(x, y) - w(x, y) = D(x, y) + a_0(x) - b_0(y)$. Choose $u \in G$ such that the homomorphism $y \mapsto D(u, y)$ is nonzero. For each x , let

$$S_x = \{y : v(x, y) = w(x, y)\}.$$

If $y \in S_x \cap S_{x+u}$, subtraction gives

$$D(u, y) + a_0(x + u) - a_0(x) = 0.$$

The solution set is empty or a coset of the kernel of a nonzero homomorphism, so it has density at most $1/2$. It follows that

$$\frac{|S_x|}{|H|} + \frac{|S_{x+u}|}{|H|} \leq 1 + \frac{|S_x \cap S_{x+u}|}{|H|} \leq \frac{3}{2}.$$

Average over x to obtain an agreement fraction at most $3/4$. For sharpness, the map xy on $\mathbb{Z}_2 \times \mathbb{Z}_2$ differs from the zero map at exactly one of the four points. $\square$

Theorem II.337.2 (Sharp robust intersection). *Let $f : G \times H \rightarrow K$, and put*

$$d_v = \text{dist}(f, \mathcal{V}), \quad d_h = \text{dist}(f, \mathcal{H}), \quad \eta = d_v + d_h.$$

If $0 \leq \eta < 1/4$, then

$$\text{dist}(f, \mathcal{A}) \leq J(\eta) := \frac{1 + \eta - \sqrt{1 - 2\eta}}{2}. \quad (\text{II.337.1})$$

There are exact equality examples for a dense set of $\eta \in (0, 1/4)$. In particular, no strictly smaller continuous upper bound as a function of η alone is valid for all finite abelian groups.

Proof. Choose nearest models $v \in \mathcal{V}$ and $w \in \mathcal{H}$. Since $\text{dist}(v, w) \leq \eta < 1/4$, Lemma II.337.1 forces their biadditive parts to agree. Subtract the common biaffine function $B(x, y) + m(x) + \ell(y)$ from f, v, w . After this subtraction the two approximants have the form $A(x)$ and $E(y)$, and their distances from the new f are still d_v, d_h .

Let p and q be the value distributions of $A(x)$ and $E(y)$. The inputs x, y are independent, so

$$\sum_z p(z)q(z) = \mathbb{P}[A(x) = E(y)] \geq 1 - \eta. \quad (\text{II.337.2})$$

Each distribution has an atom of mass at least $1 - \eta$, because their inner product is at most either largest atom. These atoms have the same location c . Indeed, if their locations were different, the agreement probability would be at most the sum of their two tail masses, hence at most $2\eta < 1 - \eta$.

Write $r = 1 - p(c)$ and $t = 1 - q(c)$. Then $0 \leq r, t \leq \eta < 1/4$, and

$$\eta \geq \mathbb{P}[A \neq E] \geq r + t - 2rt. \quad (\text{II.337.3})$$

The last inequality bounds the mass of equal nonmodal values by rt . Pointwise,

$$\mathbf{1}_{\{f \neq c\}} \leq \mathbf{1}_{\{f \neq A(x)\}} + \mathbf{1}_{\{f \neq E(y)\}} + \mathbf{1}_{\{A(x)=E(y) \neq c\}}.$$

Therefore

$$\text{dist}(f, c) \leq \eta + rt. \quad (\text{II.337.4})$$

Put $z = \sqrt{rt} < 1/2$. By the arithmetic–geometric mean inequality and (II.337.3),

$$2z - 2z^2 \leq r + t - 2rt \leq \eta.$$

The function $2z - 2z^2$ increases on $[0, 1/2]$, so

$$z \leq \frac{1 - \sqrt{1 - 2\eta}}{2}.$$

Substitution in (II.337.4) gives $\eta + z^2 \leq J(\eta)$. Adding back the common biaffine function gives (II.337.1).

For equality, take $G = H = \mathbb{Z}_n$ with n odd and $K = \mathbb{Z}_2$. Choose $S \subseteq G$ and $T \subseteq H$ with the same density u , where

$$0 < u < \frac{1 - 1/\sqrt{2}}{2}.$$

Let f be the indicator of $(S \times H) \cup (G \times T)$. There are no nonzero homomorphisms from the odd-order groups to $\mathbb{Z}_2$, so $\mathcal{V}$ consists of row functions, $\mathcal{H}$ of column functions, and $\mathcal{A}$ of constants. Optimizing a modal value in every row or column gives

$$d_v = d_h = u(1 - u), \quad \text{dist}(f, \mathcal{A}) = 2u - u^2 = J(2u(1 - u)).$$

Rational u with odd denominators are dense in the indicated interval. Their values of $\eta = 2u(1 - u)$ are dense in $(0, 1/4)$. $\square$

Remark II.337.3. The separation example in Lemma II.337.1 proves sharpness of that lemma's threshold. It does not assert that the conclusion of Theorem II.337.2 fails at precisely $\eta = 1/4$. Extending the modulus beyond its present range is a separate question. Taking $\eta = 0$ proves $\mathcal{V} \cap \mathcal{H} = \mathcal{A}$.

II.337.2 A variance bound that uses both directions

Let $\varepsilon_s^v(f)$ be the vertical arrangement failure probability from the preceding section. Define $\varepsilon_s^h(f)$ by exchanging the two input groups: choose a common horizontal increment $g \in G$, choose a uniform balanced relation among the y_i , and independently choose the x_i . The tested values are $D_g^x f(x_i, y_i) = f(x_i + g, y_i) - f(x_i, y_i)$. Thus both tests have their own balanced coordinates and their own independent labels.

Lemma II.337.4 (Two conditional collision defects). *Let $r : G \times H \rightarrow K$ satisfy $\mathbb{P}[r = 0] = 1 - \delta$. Set*

$$t_h^v = \mathbb{P}_{x,y}[D_h^y r(x, y) \neq 0], \quad t_g^h = \mathbb{P}_{x,y}[D_g^x r(x, y) \neq 0].$$

Then $0 \leq t_h^v, t_g^h \leq 2\delta$, and

$$\mathbb{E}_h t_h^v + \mathbb{E}_g t_g^h \geq 2\delta(1 - \delta). \quad (\text{II.337.5})$$

Proof. The cap follows because a nonzero difference requires at least one of its endpoints to be nonzero. For the sum, let $Z(x, y) = e_{r(x,y)}$ in the Euclidean space with orthonormal basis $(e_z)_{z \in K}$. Independence of x, y gives the orthogonal decomposition

$$Z = \mathbb{E}Z + (\mathbb{E}[Z \mid x] - \mathbb{E}Z) + (\mathbb{E}[Z \mid y] - \mathbb{E}Z) + W,$$

where W is orthogonal to all functions of x and to all functions of y . Consequently

$$\mathbb{E} \|Z - \mathbb{E}[Z \mid x]\|_2^2 + \mathbb{E} \|Z - \mathbb{E}[Z \mid y]\|_2^2 \geq \mathbb{E} \|Z - \mathbb{E}Z\|_2^2.$$

If $p_x(z) = \mathbb{P}_y[r(x, y) = z]$, the first term equals $\mathbb{E}_x(1 - \sum_z p_x(z)^2) = \mathbb{E}_h t_h^v$, since y and $y + h$ are independent uniform after averaging h . The other conditional variance is $\mathbb{E}_g t_g^h$. The total variance is $1 - \sum_z p(z)^2$, where p is the global value distribution of r . Finally,

$$\sum_z p(z)^2 \leq (1 - \delta)^2 + \delta^2,$$

because all nonzero masses sum to δ . This proves (II.337.5). $\square$

Theorem II.337.5 (Two-direction local defect profile). *For even $s \geq 4$, let $\delta = \text{dist}(f, \mathcal{A}) \leq 1/4$. Then*

$$\varepsilon_s^v(f) + \varepsilon_s^h(f) \geq P_s(\delta) := 2(1 - \delta)B_s(\delta). \quad (\text{II.337.6})$$

Proof. Subtract a nearest biaffine map and call the residual r . Each test is invariant under this subtraction, and $\mathbb{P}[r = 0] = 1 - \delta$. For $\delta = 0$ the claim is immediate. Otherwise apply Theorem II.333.2 to $D_h^y r$ with labels y , and to $D_g^x r$ with labels x . Since $E_s(t)/t$ decreases on $[0, 1/2]$ and every derivative defect is at most 2δ , the sum of the two failures is at least

$$\frac{E_s(2\delta)}{2\delta} (\mathbb{E}_h t_h^v + \mathbb{E}_g t_g^h) \geq (1 - \delta)E_s(2\delta) = P_s(\delta).$$

Here Lemma II.337.4 supplies the second inequality. The residual need not be modal in every row: only its global zero mass is used. $\square$

II.337.3 An activated inverse and optimal coefficients

Theorem II.337.6 (Local biaffine repair). *Let $s \geq 4$ be even, and set $e = \varepsilon_s^v(f) + \varepsilon_s^h(f)$. If*

$$0 \leq e \leq \frac{1}{1920s}, \quad (\text{II.337.7})$$

then

$$\text{dist}(f, \mathcal{A}) \leq \sigma_s(e), \quad (\text{II.337.8})$$

where σ_s is the inverse of P_s on $[0, 1/(16s)]$. In particular,

$$\sigma_s(e) = \frac{e}{2s} + \frac{2s-1}{4s^2}e^2 + O_s(e^3). \quad (\text{II.337.9})$$

For $s \geq 6$, the cubic coefficient is

$$\sigma_s(e) = \frac{e}{2s} + \frac{2s-1}{4s^2}e^2 + \frac{8s^2-3s-2}{12s^3}e^3 + O_s(e^4). \quad (\text{II.337.10})$$

For $s = 4$ the cubic coefficient of this upper bound is $21/128$.

Proof. By the coarse decoder in Appendix V.227 and the monotonicity $\varepsilon_4 \leq \varepsilon_s$,

$$d_v + d_h \leq 60e < \frac{1}{4}.$$

Theorem II.337.2 applies. For $0 \leq \eta < 1/4$, its proof gives $J(\eta) \leq \eta + \eta^2 \leq 2\eta$. Hence

$$\delta := \text{dist}(f, \mathcal{A}) \leq 120e \leq \frac{1}{16s}.$$

This is the required choice of inverse branch.

We verify that the branch is increasing. Write $q = 4\delta \leq 1/(4s)$. Differentiation gives

$$B'_s(\delta) = s(1-q)^{s-1} - 2(s-1)q^{s-2} + sq^{s-1}.$$

Bernoulli's inequality makes the first term at least $3s/4$. Since $s \geq 4$ and $q \leq 1/(4s)$, the absolute value of the negative term is at most $2(s-1)/(16s^2) < s/4$. Thus $B'_s \geq s/2$. Also $B_s(\delta) \leq s\delta$, by $1 - (1-4\delta)^s \leq 4s\delta$ and the nonnegative correction subtracted in B_s . Therefore

$$P'_s(\delta) = 2((1-\delta)B'_s(\delta) - B_s(\delta)) \geq s(1-3\delta) > 0. \quad (\text{II.337.11})$$

Integrating shows that $P_s(1/(16s)) > 1/(1920s)$, so the inverse exists on the whole output range in (II.337.7). Theorem II.337.5 then gives (II.337.8).

Finally, $B_s(\delta) = s\delta - 2s(s-1)\delta^2 + O_s(\delta^3)$, so

$$P_s(\delta) = 2s\delta - 2s(2s-1)\delta^2 + O_s(\delta^3).$$

For $s \geq 6$ the cubic coefficient of P_s is $\frac{4}{3}s(s-1)(4s-5)$, obtained by expanding (II.334.4). Reversal of these series gives (II.337.9) and (II.337.10). Directly using $B_4 = 4\delta - 24\delta^2 + 32\delta^3$ gives the stated exceptional cubic coefficient. $\square$

Theorem II.337.7 (Sharp coefficients through degree $s - 2$). *In the uniform expansion of the two-direction repair bound, every coefficient of σ_s through degree $s - 2$ is optimal in succession, with preceding coefficients fixed. Thus both coefficients in (II.337.9) are optimal for all $s \geq 4$, and the cubic in (II.337.10) is optimal for $s \geq 6$.*

Proof. Take $G = H = \mathbb{Z}_n$, n odd, $K = \mathbb{Z}_2$, and $f(x, y) = \mathbf{1}_{\{0\}}(x)$. There are no nonzero homomorphisms from either domain group to K , so the biaffine maps are constants and

$$\delta = \text{dist}(f, \mathcal{A}) = \frac{1}{n}.$$

All vertical tests pass. A horizontal increment $g = 0$ also always passes. For $g \neq 0$, the derivative is supported on exactly two x -values. The x_i in the horizontal test are independent, so its s tested derivatives are independent Bernoulli random variables with parameter $2/n$ in $\mathbb{Z}_2$. The balanced condition constrains only their y -coordinates. Thus the failure is odd parity, and exactly

$$e = C_s(\delta) := \frac{1 - \delta}{2} (1 - (1 - 4\delta)^s). \quad (\text{II.337.12})$$

Comparison with P_s gives

$$C_s(\delta) - P_s(\delta) = 2^{2s-2} \delta^{s-1} (1 - \delta)(1 - 2\delta). \quad (\text{II.337.13})$$

Both functions have derivative $2s$ at zero. Their inverse power series therefore agree through degree $s - 2$. Letting odd $n \rightarrow \infty$ gives examples with $e \rightarrow 0$ on which each of those coefficients is attained. A reduction of the first differing coefficient, with earlier ones fixed, would contradict this sequence. $\square$

Remark II.337.8. The factor in (II.337.13) begins at degree $s - 1$. The proof therefore does not establish optimality at degree $s - 1$ or beyond. In particular, the cubic $21/128$ at $s = 4$ is a valid upper-bound coefficient, not a claimed sharp one. Nor is P_s asserted to be the full finite-error envelope.

II.337.4 How often should each direction be tested?

Corollary II.337.9 (Optimal fixed allocation at leading order). *Suppose a tester chooses the vertical arrangement test with probability $p \in (0, 1)$ and the horizontal test with probability $1 - p$. Write*

$$\varepsilon_{\text{mix}} = p\varepsilon_s^v + (1 - p)\varepsilon_s^h, \quad a = \min\{p, 1 - p\}.$$

If $\varepsilon_{\text{mix}} \leq a/(1920s)$, then

$$\text{dist}(f, \mathcal{A}) \leq \sigma_s(\varepsilon_{\text{mix}}/a) = \frac{\varepsilon_{\text{mix}}}{2sa} + O_{s,p}(\varepsilon_{\text{mix}}^2).$$

The leading coefficient $1/(2sa)$ is optimal for each fixed p . It is minimized by $p = 1/2$, where it equals $1/s$.

Proof. Since $\varepsilon_{\text{mix}} \geq a(\varepsilon_s^v + \varepsilon_s^h)$, Theorem II.337.6 applies. The sparse-row example of Theorem II.337.7, or its transpose, puts all failure in the less frequently tested direction. Its error is $\varepsilon_{\text{mix}} = aC_s(\delta) = 2sa\delta + O_s(\delta^2)$, proving the lower obstruction. Maximizing $\min(p, 1 - p)$ gives the stated allocation. $\square$

Status. This section resolves the local full-product two-direction repair problem with sharp first-order constants and more precise asymptotics. Sparse domains, optimal finite-error envelopes, and stability of the robust-intersection equality examples remain open.

II.338 Source 67's Section 5: the interpolation count at Lemma 16.10 [write]

Added in the eleventh write, 7 October 2026.

The next section is source 67's Section 5. *Logically it continues source 39's chapter* on shared anchors at Lemma 16.10 (Section II.151) and source 19's sampling sections; it is printed here because no number printed earlier may move.

What is proved. If anchor i carries a list of L_i base functions and fibre classes have degree below t , then for $1 \leq t \leq r$ every certified point (sampled, or with t anchors in its class) is covered by at most $E_t(L_1, \dots, L_r)$ functions (Theorem II.339.1): the ordinary t -anchor interpolants also cover the sampled points, so source 39's count $\sum_i L_i + E_t$ loses its first term (for affine fibres, $rL + \binom{r}{2}L^2$ becomes $\binom{r}{2}L^2$). The complete formula, with empty lists, is Theorem II.339.2. The count is exact for the local interface: dense examples over large prime fields force all E_t candidates even against candidates of arbitrary fixed degree in the base variable (Theorem II.339.6), with a product lower bound on the uncovered mass, so that $M \geq (1 - \delta - o(1))E_t$ when a fraction δ may be deleted (Corollary II.339.8), and a cell-count trade-off (Corollary II.339.10). The elementary symmetric counting is credited to source 39, as source 67 says. The closing comparison of Lemma 16.10 is not addressed.

Relation to the formal statements (checked at HEAD fb9602e55). `lemma_16_10` has changed meaning since source 67's arrival. The catalogued predicate that source 50 refuted (Section II.226) is now `lemma_16_10_printed_unit_encoding` (Section16.lean, line 852, unchanged in content; its negation `lemma_16_10_packaged_premises_counterexample` now targets that name), and `lemma_16_10` (line 912) states a repaired lemma: from the same two packaged premises, for every loss $\rho \in (0, 1]$, a cover of ϕ_1 on B_1 off $\rho|P|$ points by at most $\max(b, \binom{r}{2}b^2)$ multilinear graphs on a partition of P into proper boxes, where $r = \lceil 6 \max(1, q_r)/\sigma \rceil$ anchor slices are sampled at $\sigma = \rho/4$; the graph count has the form of the compressed affine count below. Its companion `lemma_16_10_holds` (Proofs16Lemma10.lean, line 12; 456b6e3ad, 7 October 2026, 12:20) is counted in the ledger (114 exact companions, 6 open statements) but is *not yet kernel-checked*: `FORMALIZATION_STATUS.txt` records that for the repair only `Section16.lean` and 61 of its upstream modules were compiled, and that full verification awaits the next complete audit. The repaired form does not supply the unit-parameter conclusion that the printed proof of Theorem 16.2 uses, and Theorem 16.2 stays open. *The affine case of Theorem II.339.1 has been formalized since source 67's arrival*, in the repository's own form: `Section16SampledOrAnchoredOn.compressed_multilinear_cover` (Proofs16CompressedInterpolation.lean, line 56; 081abc4d3, 7 October 2026), a family of $\max(p, \binom{r}{2}p^2)$ candidates for arbitrary samples with replacement, extending source 39's interface `Section16SampledOrAnchoredOn` (Proofs16ShortColumns.lean, line 27), and carried through the repository's Section 16 affine lift (`section16UniformLiftGraphBudget_lt_previous`, `Proofs16CompressedLiftImprovement.lean`, line 39). `FORMALIZATION_STATUS.txt` names this section as its origin. The higher-degree count, the lower constructions and the deletion trade-offs are not formalized.

II.339 The exact number of interpolation candidates

Source 67, delivered Section 5.

The shared-anchor argument in Gowers's Lemma 16.10 [1] has two logically separate costs: finding anchors that certify most points, and converting the slice lists at those anchors into a list on the whole base cell. Sources 19 and 39 in the existing ProveIt report [269] substantially improve the first cost and preserve the labels of the second. Source 39's Section 10 gives, for degree less than t , the candidate bound

$$\sum_{i=1}^r L_i + \sum_{\substack{I \subseteq [r] \\ |I|=t}} \prod_{i \in I} L_i,$$

and its further-research section asks whether algebraic coincidences can compress these lists and whether a matching lower-bound model exists. Here the direct-anchor term is removed, and the remaining elementary symmetric polynomial is shown to be the exact universal answer for the fixed-cell interface. A quantitative lower bound continues to hold after deletion of points, and gives a tradeoff between the number of further cells and their list sizes.

This is a local assertion with a concrete optimality scope. The construction satisfies the same slice-list and polynomial-fibre hypotheses, with one class in every fibre and a domain of density arbitrarily close to one. It is not asserted to satisfy every additional product-property hypothesis of the full Gowers induction. The ordinary interpolation count is closely related to polynomial list recovery at the minimum agreement threshold; no priority claim for Lagrange interpolation or for the general coding-theoretic viewpoint is intended.

II.339.1 The local interface and the smaller upper bound

Let $\mathbb{F}$ be a field, let $J \subseteq \mathbb{F}$ be finite, and fix distinct anchors $a_1, \dots, a_r \in J$. Write $R = \{a_1, \dots, a_r\}$. Let H be a finite base set and let $V \subseteq \mathbb{F}^H$ be a linear space. A partial function

$$\phi : B \longrightarrow \mathbb{F}, \quad B \subseteq H \times J,$$

has fibre $B_h = \{y : (h, y) \in B\}$. In every fibre suppose that B_h is partitioned into classes, and on each class C the function $y \mapsto \phi(h, y)$ agrees with a polynomial $P_{h,C}$ of degree less than t . The number of classes will not enter the candidate theorem.

At each anchor a_i suppose that a nonempty list

$$\mu_{i,1}, \dots, \mu_{i,L_i} \in V, \quad L_i \geq 1,$$

covers the partial slice: whenever $(h, a_i) \in B$, its value is $\mu_{i,j}(h)$ for at least one j . A common good subset of H can simply be substituted for H ; the theorem makes no additional deletion.

A point of B is *certified* if it is itself sampled, or if its class contains at least t distinct anchors. This is the certification rule of source 39: sampled points are retained even when their class has fewer than t anchors. Put

$$E_t(L_1, \dots, L_r) := \sum_{\substack{I \subseteq [r] \\ |I|=t}} \prod_{i \in I} L_i.$$

Theorem II.339.1 (No separate direct-anchor candidates). *Suppose $1 \leq t \leq r$. All certified points are covered by at most*

$$T = E_t(L_1, \dots, L_r) \tag{II.339.1}$$

functions in

$$V \otimes \text{Poly}_{<t}(J) = \left\{ \sum_{j=0}^{t-1} v_j(h) y^j : v_j \in V \right\}.$$

In particular, if every $L_i \leq L$, the count is at most $\binom{r}{t} L^t$. For affine fibres and separately affine base lists, the count is $\binom{r}{2} L^2$, and every candidate is separately affine in all base variables and in y .

Proof. For a t -subset $I \subseteq [r]$ and one index $j_i \in [L_i]$ for each $i \in I$, form the Lagrange interpolant

$$F_{I,(j_i)}(h, y) = \sum_{i \in I} \mu_{i,j_i}(h) \prod_{\substack{k \in I \\ k \neq i}} \frac{y - a_k}{a_i - a_k}. \quad (\text{II.339.2})$$

All denominators are nonzero and all coefficients in the y -expansion belong to V . There are T indexed choices, possibly with repetitions.

For a nonsampled certified point, choose t anchors from its class and a correct list member at each of them. The two polynomials in y , the class polynomial and the candidate evaluated at h , have degree less than t and agree at t distinct points. They therefore coincide.

For a sampled point (h, a_i) , choose *any* t -subset I containing i , select a correct list member at i , and select arbitrary members of the other nonempty lists. At $y = a_i$, all other Lagrange coefficients vanish and the i th coefficient is one. Thus this candidate covers the sampled point. Values at the other chosen anchors need not belong to the same class, or even to the domain of ϕ ; their candidate functions are nevertheless defined on the whole base cell. This is precisely why the direct-anchor candidates are redundant. $\square$

[write] **7 October 2026; sources 60–69.** *Formalized since its arrival*, in the affine case and in the repository's own form (checked at HEAD fb9602e55): `Section16SampledOrAnchoredOn.compressed_multilinear_cover` (`Proofs16CompressedInterpolation.lean`, line 56; 081abc4d3) gives a family of $\max(p, \binom{r}{2} p^2)$ candidates for arbitrary samples with replacement, and `section16_compressed_budget_lt_old` (`Proofs16CompressedCandidateBounds.lean`, line 42) the strict comparison with $rb + r^2 b^2$; `FORMALIZATION_STATUS.txt` cites this theorem. The case $t \geq 3$ and the lower constructions are not formalized. Source 39's count is its Theorem II.160.1 and Corollary II.160.2.

Theorem II.339.2 (All anchor counts, including empty lists). *Allow $r \geq 0$, $t \geq 1$, and $L_i \geq 0$. Let $A = \{i : L_i > 0\}$ and $a = |A|$. With the same certification and slice-cover hypotheses, every certified point is covered by at most*

$$\mathcal{L}_t(L_1, \dots, L_r) = \begin{cases} 0, & a = 0, \\ \max_i L_i, & 0 < a < t, \\ E_t(L_1, \dots, L_r), & a \geq t \end{cases} \quad (\text{II.339.3})$$

candidates in $V \otimes \text{Poly}_{<t}(J)$. For $t = 1$ this becomes $\sum_i L_i$, including the all-empty case.

Proof. An anchor with an empty covering list cannot meet the domain B at any base point under consideration. Removing those anchors changes neither the sampled domain points nor the number of anchors in a class. If none remain, there are no certified points. If at least t remain, apply Theorem II.339.1; terms involving an empty list vanish from E_t .

If $0 < a < t$, no class can meet t anchors, so all certified points are sampled. Put $L_* = \max_i L_i$. Pad every nonempty list by repeating its last member until it has length L_* . For each $1 \leq j \leq L_*$,

interpolate the j th members across all a active anchors. The resulting polynomial in y has degree at most $a - 1 < t$, with coefficients in V . At each active anchor the L_* candidates include every member of that anchor's list. They therefore cover all sampled points. $\square$

The change at $a = t$ reflects the certification rule. Below the threshold, only anchor values must be covered, and their lists can be aligned. At the threshold, the rule certifies nonsampled values, and all independent interpolation choices may become necessary.

II.339.2 Generic finite-field lists

The lower bound requires genuinely different interpolants. The next elementary lemma supplies them over an explicit range of prime fields. For $u > r$, use the convention $E_u(L_1, \dots, L_r) = 0$.

Lemma II.339.3 (Lists with no interpolation collisions). *Fix integers $r \geq t \geq 1$ and $L_1, \dots, L_r \geq 1$, and set*

$$D = L_{i_*} - 1 + E_t((L_i)_{i \neq i_*}), \quad L_{i_*} = \max_i L_i. \quad (\text{II.339.4})$$

Let p be a prime with $p > \max\{r, D\}$. For any distinct $a_1, \dots, a_r \in \mathbb{F}_p$, one can choose sets $A_i \subseteq \mathbb{F}_p$ of cardinality L_i such that the following hold:

1. *no polynomial of degree less than t takes values in A_i at more than t of the anchors;*
2. *the interpolants obtained by choosing a t -subset I and one value from each A_i , $i \in I$, are all distinct. There are exactly $T = E_t(L_1, \dots, L_r)$ of them.*

Proof. Process the anchors in an order for which $L_1 \leq \dots \leq L_r$, after relabelling if needed. Suppose the lists at the previous anchors have been chosen. At anchor a_i , forbid the values at a_i of all polynomials interpolated from t of the previous lists. There are at most $E_t(L_1, \dots, L_{i-1})$ forbidden values. Since

$$L_i + E_t(L_1, \dots, L_{i-1}) \leq L_r + E_t(L_1, \dots, L_{r-1}) = D + 1 \leq p,$$

one can choose L_i distinct allowed values. This gives a deterministic construction: for example, choose the first available residues in increasing order.

If a polynomial of degree less than t meets the final lists at $t + 1$ distinct anchors, consider the last of those anchors in the construction. Its value was forbidden by the preceding t anchors, a contradiction. This proves the first conclusion. For the second, suppose two indexed interpolants coincide. If their anchor sets coincide, their values at each anchor force their ordinate selections to coincide. If their anchor sets differ, their common polynomial meets the lists at every anchor in their union, which has at least $t + 1$ members. This contradicts the first conclusion. $\square$

Remark II.339.4 (Meaning of the generic condition). For $t = 2$, the condition says that no nonvertical line through chosen points on two anchor lines passes through a chosen point on a third anchor line. Thus all cross-anchor secants are different. The lemma gives a deterministic finite construction, not an unspecified asymptotic genericity assumption. It is not intended to optimize the true minimum field size. An independent union-bound proof would give the sufficient condition $p > \sum_i \binom{L_i}{2} + E_{t+1}(L_1, \dots, L_r)$, usually larger.

Proposition II.339.5 (The best order for the greedy cardinality bound). *Among all anchor orders, the maximum of the stepwise upper bounds*

$$L_i - 1 + E_t(L_1, \dots, L_{i-1})$$

is minimized by a nondecreasing order of the list sizes, and its value there is D from (II.339.4).

Proof. Consider adjacent sizes $x \leq y$ after a fixed prefix, and put $A = E_t(\text{prefix})$, $B = E_{t-1}(\text{prefix})$, with $E_0 = 1$. The maximum of their two costs in the order x, y is $A + y + xB - 1$. In the order y, x it is $A + \max\{y, x + yB\} - 1$. If $B = 0$ the maxima coincide. If $B \geq 1$, $y + xB \leq x + yB$ because $(y - x)(B - 1) \geq 0$. Thus removing an inversion cannot increase the maximum. All later prefix elementary symmetric polynomials are unchanged. Repeated adjacent swaps prove optimality. For a nondecreasing order the stepwise bounds are nondecreasing, so the last equals the maximum, namely D . $\square$

II.339.3 Exact optimality and deletion robustness

Write $x_+ = \max\{x, 0\}$. A polynomial has bidegree at most $(d, t - 1)$ if it has degree at most d in the base variable h and degree at most $t - 1$ in the fibre variable y , separately.

Theorem II.339.6 (Dense sharp examples and an explicit uncovered-mass bound). *Fix $r \geq t \geq 1$, $L_1, \dots, L_r \geq 1$, $d \geq 1$, and $n \geq 1$. Let T and D be defined in (II.339.1) and (II.339.4), and suppose that p is prime with*

$$p > \max\{r, D, Tn - 1\}. \quad (\text{II.339.5})$$

There exist an interval $H \subseteq \mathbb{F}_p$ of size Tn , r anchors in $J = \mathbb{F}_p$, and a partial function $\phi : B \rightarrow \mathbb{F}_p$ with $B \subseteq H \times J$ such that:

1. *every fibre has one polynomial class of degree less than t ;*
2. *the slice at anchor i is covered by L_i constant functions;*
3. *every point of B is certified by the r anchors;*
4. *the domain density is*

$$\frac{|B|}{|H|p} = 1 - \frac{r - t}{p}; \quad (\text{II.339.6})$$

5. *for every list of M polynomials of bidegree at most $(d, t - 1)$, the number U of domain points covered by none of them satisfies*

$$\boxed{U \geq (T - M)_+(n - dM)_+(p - r - (t - 1)M)_+}. \quad (\text{II.339.7})$$

Consequently, if

$$n > d(T - 1), \quad p - r > (t - 1)(T - 1), \quad (\text{II.339.8})$$

every exact list cover requires at least T candidates. Exactly T are sufficient. In particular, the universal bound $\binom{r}{2}L^2$ for affine fibres is best possible already for one base variable, even if the competing candidates may have any fixed base degree d .

Proof. Take the lists A_i from Lemma II.339.3, and enumerate their T distinct interpolants as $P_1, \dots, P_T$. Let I_ν be the t anchors used to define P_ν . Set

$$H = \{0, 1, \dots, Tn - 1\} \subseteq \mathbb{F}_p, \quad H_\nu = \{(\nu - 1)n, \dots, \nu n - 1\}.$$

The size hypothesis makes these disjoint n -element blocks. Let $S = \mathbb{F}_p \setminus R$, so $|S| = p - r$. For $h \in H_\nu$ put

$$B_h = S \cup \{a_i : i \in I_\nu\}, \quad \phi(h, y) = P_\nu(y) \quad (y \in B_h). \quad (\text{II.339.9})$$

The whole fibre is one polynomial class and contains exactly t anchors. Thus it is entirely certified. At an anchor a_i the only defined values are elements of A_i , so A_i supplies a list of L_i constant functions. Every fibre has size $p - r + t$, proving (II.339.6).

Let the proposed cover be $F_1(h, y), \dots, F_M(h, y)$. A candidate can coincide as a polynomial with at most one of the distinct targets $P_\nu(y)$. Hence at least $(T - M)_+$ targets have no identical candidate. Fix such a target.

For each j , the difference

$$F_j(h, y) - P_\nu(y) = \sum_{b=0}^{t-1} c_{j,b}(h) y^b$$

is a nonzero polynomial, and some coefficient $c_{j,b}$ is a nonzero polynomial of degree at most d . There are at most d values of h at which every coefficient vanishes, that is, at which $F_j(h, \cdot)$ and P_ν agree identically in y . Across all M candidates, at most dM rows have such a full-row coincidence. On each of the remaining at least $(n - dM)_+$ rows of H_ν , each candidate differs from P_ν by a nonzero polynomial of degree at most $t - 1$ in y . It agrees at at most $t - 1$ points of S . Their union therefore covers at most $(t - 1)M$ of the $p - r$ nonsampled columns on that row. At least $(p - r - (t - 1)M)_+$ points remain uncovered. Sum over the unmatched targets. All counted points belong to $H \times S \subseteq B$, which proves (II.339.7).

If $M < T$, conditions (II.339.8) make every factor strictly positive, so exact coverage is impossible. Conversely, the list $P_1(y), \dots, P_T(y)$ covers B and has the required bidegree. $\square$

Corollary II.339.7 (A single explicit sufficient prime condition). *With T, D, r, t, d fixed, put $n = d(T - 1) + 1$. Every prime*

$$p > \max\{D, T[d(T - 1) + 1] - 1, r + (t - 1)(T - 1)\}$$

admits an example with exact minimum T candidates. For the affine case and separately affine candidates, one may take $t = 2$, $d = 1$, $n = T$, and any prime

$$p > \max \left\{ L_{i_*} - 1 + \sum_{\substack{i < j \\ i, j \neq i_*}} L_i L_j, T^2 - 1, r + T - 1 \right\}.$$

Proof. The displayed inequalities give (II.339.5) and (II.339.8). In the general formula the last term is at least r , as required by Lemma II.339.3. $\square$

Corollary II.339.8 (Asymptotically exact cost with permitted deletion). *Fix $r, t, L_1, \dots, L_r, d$ and $0 \leq \varepsilon < 1$. There are examples of the preceding type with domain density tending to one such that every list of $M < T$ bidegree- $(d, t - 1)$ candidates leaves an ambient fraction at least*

$$\left(1 - \frac{M}{T}\right) \left(1 - \frac{dM}{n}\right) \left(1 - \frac{r + (t - 1)M}{p}\right) \quad (\text{II.339.10})$$

uncovered, whenever the displayed factors are nonnegative. In particular, for $n \rightarrow \infty$ and primes $p \geq Tn$ tending to infinity, an ambient deletion allowance of ε forces

$$M \geq (1 - \varepsilon - o(1))T.$$

Conversely, retaining any M of the target polynomials covers all points in the corresponding M row blocks and possibly further points, so deleting at most a fraction $1 - M/T$ of B always suffices. Thus the leading deletion-list tradeoff is sharp.

Proof. Divide (II.339.7) by $|H|p = Tnp$. The two last factors tend to one uniformly over the finite range $0 \leq M < T$. The upper comparison deletes all unselected blocks; these blocks have the same domain mass because every row has $p - r + t$ points. $\square$

Proposition II.339.9 (Sharpness below the anchor threshold). *Suppose that exactly a lists are nonempty, where $0 < a < t$, and put $L_* = \max_i L_i$. Fix $d \geq 1$ and choose $n > d(L_* - 1)$. Over every sufficiently large prime field, there is an example satisfying the slice-list and certification hypotheses whose exact minimum number of bidegree- $(d, t - 1)$ candidates is L_* . Empty lists and the all-empty case are therefore handled by the sharp piecewise budget in (II.339.3).*

Proof. It suffices to take a prime $p > \max\{r, L_*, L_*n - 1\}$ and r distinct anchors. Partition the interval of L_*n base points into blocks $H_0, \dots, H_{L_*-1}$ of size n . For every active anchor i , use the list of constant functions $0, 1, \dots, L_i - 1$. At a row $h \in H_\nu$, define the domain to consist of just the a active anchors and put

$$\phi(h, a_i) = \min\{\nu, L_i - 1\}.$$

These a values lie on one polynomial of degree at most $a - 1 < t$, so the whole fibre can again be one polynomial class. Every domain point is sampled. The interpolation construction in Theorem II.339.2 supplies L_* candidates.

Choose an anchor i_* with $L_{i_*} = L_*$. At this column, the values on the successive blocks are the L_* distinct constants $0, \dots, L_* - 1$. For any $M < L_*$ proposed candidates, at least $L_* - M$ of these constants are not identically equal to any of the univariate polynomials $h \mapsto F_j(h, a_{i_*})$. On the block for such a constant, each difference is a nonzero polynomial of degree at most d , and the union of the M candidates matches at at most $dM < n$ rows. Exact coverage is impossible. When all lists are empty there are no certified points, so the correct minimum is zero. $\square$

II.339.4 What additional partitioning can and cannot save

The fixed-cell lower bound should not be confused with a prohibition on making more cells. In fact the construction explicitly shows the natural exchange: the T consecutive base blocks can each be assigned one candidate. The correct invariant of unrestricted further partitioning is total candidate occurrences.

Corollary II.339.10 (Cell-count versus list-size tradeoff). *Use an example from Theorem II.339.6. Partition its base rectangle into finitely many cells, and on cell j permit a list of M_j functions that are restrictions of ambient polynomials of bidegree at most $(d, t - 1)$. Put $M = \sum_j M_j$. The uncovered-mass bound (II.339.7) still holds with this M . Under (II.339.8), an exact cover requires*

$$\sum_j M_j \geq T. \tag{II.339.11}$$

Consequently, b cells with at most Q candidates each require $bQ \geq T$. This product bound is attainable whenever the T target blocks can be grouped into b groups of size at most Q ; consecutive groups give ordinary rectangular arithmetic-progression cells.

Proof. Choose an ambient polynomial extension of every local candidate. Discarding the cell restrictions can only increase coverage. This gives a global cover with at most M candidates, to which Theorem II.339.6 applies. For attainment, group consecutive blocks H_ν , retain all columns, and list the target polynomials belonging to each group. Every graph point in that cell is covered. $\square$

Remark II.339.11 (Higher base dimensions). For the separately affine case, take a Cartesian product of the example with any finite sets of additional base coordinates and let ϕ ignore those coordinates. For every fixed choice of their values, each competing separately affine candidate becomes bidegree $(1, t-1)$ in the original base and fibre variables. Applying the same bound on every slice proves identical proportional lower bounds in any positive base dimension. This does not impose an exponential penalty in the number of added coordinates.

II.339.5 Consequences for the Section 16 interface

For affine fibres, source 39's labelled refinement gives a common per-anchor list bound

$$L = Q\left(\frac{\epsilon}{rb}, \gamma, k\right)^b,$$

where b is the inherited slice complexity and the base scale is controlled separately. At every final base cell with $r \geq 2$, the new upper theorem replaces the old expression

$$rL + \binom{r}{2}L^2$$

by $\binom{r}{2}L^2$, with exactly the same anchors, base partition, exceptional set, and admissibility hypotheses. Its proof uses no geometric synchronization step. If a later synchronization only restricts already constructed candidates, the improvement survives unchanged.

The exact lower bound answers source 39's question about a generic matching lower-bound model: the remaining elementary-symmetric cost cannot universally be lowered from only those slice-list and certification hypotheses. This remains true with one polynomial class per fibre, nearly full domain, arbitrary fixed extra degree in the base variable, and a small allowed deletion fraction. Compression therefore needs a stronger premise, for example more than t compatible anchors per target polynomial, relations tying choices at different anchors, or information from the product property. The example identifies which information the present interface has discarded.

The theorem does not close Gowers's numerical recurrence. The source controls still require an independent audit of scale, outside error, and list size through Lemmas 16.9–16.11. The result settles the local candidate optimum and states the precise partition tradeoff; it makes no assertion of a new bound for $r_k(N)$ or for van der Waerden numbers.

II.339.6 Further questions opened by the sharp interface

1. Suppose each target class has at least $u > t$ valid anchors and each such anchor is covered by its labelled list. Determine a sharp candidate bound that uses this surplus agreement. A first comparison is $E_t(L)/\binom{u}{t}$ when a globally compatible polynomial is already specified by its anchor values, but a pointwise base-dependent choice must not be silently replaced by such global compatibility.

2. In Gowers's full product-property setting, can the generic collision-free examples be excluded quantitatively? A positive result should provide a verifiable relation among the anchor lists and a concrete improvement over the exact fixed-interface optimum.
3. For prescribed bounded prime fields, determine the maximum number of distinct t -anchor interpolants when the sufficient genericity condition $p > D$ is unavailable. This is a finite incidence/list-recovery problem rather than an asymptotic permission to choose an arbitrarily large field.
4. Optimize the finite deletion bound (II.339.7). Its asymptotic leading term is sharp, but the factors dM and $(t-1)M$ are union bounds; incidences shared between candidates can improve the finite correction.
5. Impose a lower width on every further arithmetic cell, rather than merely charging total candidate occurrences. Determine the best tradeoff between width, number of cells, deletion, and local list size. The elementary construction here permits T consecutive blocks of width n , so any stronger lower bound needs a row-label placement that prevents such inexpensive separation.

Part III

Cube and progression counts

Contents. Counting consequences of uniformity. Sections III.1–III.2: source 03, delivered Sections 5–6 (progression counting by a density-weighted telescope; the exact fourth-order term of the centered cube count on every finite abelian group, its torsion correction, sharpness for sets, higher coefficients). Sections III.3–III.4: source 04, delivered Sections 6–7 (the sharp real centered gap $\|f\|_{U^3}^8 \geq 2 \|f\|_{U^2}^8$ on odd-order groups, the sharp leading constant $6\sqrt{2}$ of the three-cube bound, the fractional correction exponent $16/3$, and density-weighted progression counting with the lifting lemma).

Overlaps (dated notes). Cancellation through degree three and the quartic expansion (03 on every group, 04 on odd-order groups; $P_d = R_d$); the $\mathbb{Z}/3\mathbb{Z}$ cosine polynomial (identical in both); the generalized von Neumann bound and the telescoping theorem (independently written in both). **Cross-source answers.** Source 04’s Theorem III.3.4 answers source 03’s Research question V.14.1 for $d = 3$, with $c_3^{\text{odd}} = 2^{-1/2}$. Source 04’s Lemma III.4.5 is the only statement of the five sources that is formalized (`hasNatAP_of_short_modular_sequence`).

Reading conventions of this Part (the full dictionary is Section F.3).

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
cube count	03: $Q_d(F)$, $m = 2^d$ vertices, $\eta = \ f\ _{U^d}$; 04: $Q_d(W)$, $q = 2^d$, $u = \ f\ _{U^d}$, $v = \ f\ _{U^2}$. <i>04’s q here is not its kernel degree q of Part II.</i>
R_d, P_d, T_d	03: R_d = number of ordinary parallelograms, T_d = binary tetrahedra, $K_d = R_d + T_d$, $\mathcal{R}_d$ the remainder; 04: P_d = number of parallelograms (= 03’s R_d), R_d, R_6, R_7 remainders. <i>04’s R_d is not 03’s R_d.</i>
$\mathcal{E}, B$	03: $\mathcal{E}(f) = \ f\ _{U^2}^4$, $\mathcal{E}_{(2)}(f)$ the order-two (quotient) energy, $B(S)$ the telescoping recursion; 04: $\mathcal{E}(\lambda)$ the additive energy of $\lambda_r = a_r ^2$, $B_5(f)$ the quintic correlation, $D(S)$ the recursion (= 03’s $B(S)$).
slopes, α	03: slopes a_i with invertible differences, Λ_a , $\eta = \ 1_A - \delta\ _{U^{k-1}}$; 04: admissible slopes c_i , Λ_c , u . Both: Gowers’s α -uniformity means $\ f\ _{U^d}^{2^d} \leq \alpha$; <i>α is a norm power, not a norm.</i>

Added in the second write (6 October 2026): sources 06–10. Contents. Section III.5: source 06, delivered Section 3 (a generalized von Neumann bound keeping two distinguished functions, first and second telescoping, a progression discrepancy quadratic in the uniformity norm, the sharp three-term variance factor, an explicit modular existence criterion). Sections III.6–III.15: source 08, delivered Sections 2–10 and Appendix A (cancellation through degree three, the exact quartic term with its order-two correction, the complete odd-order three-cube polynomial, the counterexample to the quartic lower bound with its exact certificate, sharpness for sets, Fourier stability, inhomogeneous weights, the character-kernel survival criterion). Sections III.16–III.27: source 09, delivered Sections 2–10, 13 and Appendices A–B, the fullest treatment of the cube count (complete four- and five-vertex classifications on every finite abelian group with Stirling multiplicities, the exact fourth and fifth orders, signs and sharpness through order six, the determinant-five configuration and the all-prime family, heterogeneous expansions, and the relation to progression bounds). Sections III.28, III.29 and III.30: source 10, delivered Sections 5, 6 and 9 (a torsion-sensitive mixed- L^p generalized von Neumann inequality and its sharpness; density-weighted counting and modular criteria; the exact dimension threshold for polynomial colourings).

Overlaps (dated notes at both ends). Cancellation through degree three: 03, 04, 08, 09. The exact quartic term on every group: 03, 08, 09 ($R_d = P_d$, T_d , $K_d = F_d = A_d$). The complete odd-order three-cube polynomial: 08’s Theorem III.9.2 contains 04’s support count and is the $d = 3$ case of 09’s Theorem III.21.1. The quartic lower bound with coefficient 12 is false: 03 (on $\mathbb{Z}/3\mathbb{Z}$), 08 (on $\mathbb{Z}/11\mathbb{Z}$, with the explicit ratio $15337/1296$) and 09 (on primes $q \geq 17$) show it; Remark W1 keeps it on record with the counterexample.

Sharpness for sets by rounding: 03, 04, 08, 09. Progression counting: the linear telescope of 03 and 04 is also 06's Theorem III.5.2 and 10's Theorem III.29.1; 06 adds a quadratic bound, 10 the torsion kernels.

Reading conventions for the sections of sources 06–10 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
cube count	08: $\mathcal{C}_d(F)$, $p = 2^d$, $t = \ f\ _{U^d}$, $\mathcal{D}_d(F) = \mathcal{C}_d(F) - \delta^p$; 09: $\mathcal{C}_d(F)$, $m = 2^d$, $\varepsilon = \ f\ _{U^d}$, $a_\chi = \widehat{f}(\chi)$. Both are 03's and 04's Q_d . 08's p is not a prime.
$R_d, P_d, T_d, F_d, A_d, K_d$	Parallelograms: 08's $R_d = 09$'s $P_d = 03$'s $R_d = 04$'s P_d . Non-parallelogram binary planes: 08's $T_d = 09$'s $T_d = 03$'s T_d . All binary planes: 08's $F_d = 09$'s $A_d = 03$'s K_d . 08's $K_d = \sum_{r \geq 4} \binom{p}{r}$ is not 03's K_d ; 08's F_d is not the floor count F_k of Part IV.
Q_d, S_d	09: Q_d the number of star supports, S_d of ternary supports. 08: $Q_d(f) = R_d E_4(f) + T_d E_{\langle 2 \rangle}(f)$, the quartic coefficient. <i>Neither is the cube count Q_d of 03 and 04.</i>
Fourier statistics	$E_4(f)$ (08) = $E_4(f)$ (09) = $\mathcal{E}(f)$ (03) = $\ f\ _{U^2}^4$; $E_{\langle 2 \rangle}(f)$ (08) = $T_2(f)$ (09) = $\mathcal{E}_{\langle 2 \rangle}(f)$ (03); $S_5(f)$ (08) = $M_5(f)$ (09) = $B_5(f)$ (04); $T_{3,5}(f)$ (09) = $M_{S_3}(f)$ (03); κ_d^{odd} (08) = c_d^{odd} (03); the subset moments $\Lambda_S(f)$ (08) = $T_S(f)$ (09) = $M_S(f)$ (03).
progressions	06: Λ_a , $\varepsilon = \ 1_A - \delta\ _{U^{k-1}}$ (a norm), $S_k(\delta)$, $B_k(\delta) = \sum_t (t+1)\delta^t$, E_k . 10: $\kappa_G(c) = G[c] $, the kernel product K , $\tau_j(G)$, B_j , D_j , $T_{k,G}(\delta)$, $S_k(\delta)$ (= 06's). 06's $B_k(\delta)$ is not the translation budget B_k of Part IV; 10's κ_G is none of the other κ 's.
colourings, $h(\ell)$	10, Section III.30: $k = 2m$, a prime $p > k$, a colouring Q with degrees D_ℓ , $s_\ell = \lceil D_\ell/2 \rceil$, $B = \sum_\ell s_\ell^2$, the field norm N_t (<i>this m is $k/2$</i>). 09: $h(\ell)$ the least support with isolated ℓ -torsion, $\mathcal{L}_S$ the row lattice, $K_G(S)$ the character constraint group.

Added in the third write (6 October 2026): sources 11–14. Contents. Sections III.31–III.33: source 13, delivered Sections 7–9 (centered cube counts over $\mathbb{F}_p^n$: coloop cancellation, the exact quartic coefficient, the three-cube table and an exact spectral identity with orbit moments and a spectral deficit; interval discrepancy with the sharp exponent 1/4 for Lemma 3.5; telescoping progression counts). Sections III.34–III.38: source 14, delivered Sections 2, 3, 6, 7 and Appendix A (the exact secondary coefficient $(9/4)^{1/3}$ of the three-cube extremal problem, two-harmonic rigidity and indicator sharpness; exponent six on groups of exponent five; exact supersaturation for polynomial colourings; quantitative examples; the support certificate).

Overlaps (dated notes). The three-cube identity: source 08 on every odd-order group, source 13 over $\mathbb{F}_p^n$, source 09 for every d ; the table $(1, 0, 0, 0, 12, 8, 28, 8, 1)$: 04, 08, 13, 14; the quartic coefficient: 03, 08, 09, 13; the telescope: 03, 04, 06, 10, 13; the colouring threshold: 10, reproved by 14. **Cross-source answers.** Source 04's Question 1 is *settled* by source 14 (Remark W4): the secondary coefficient exists and equals $(9/4)^{1/3}$. Source 04's Question 2 is answered in U^2 and fails in L^2 (source 14's Theorem III.34.3 and the remark after it).

Reading conventions for the sections of sources 11–14 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
cube count	13: $\mathcal{Q}_d(\delta + f)$, $m = 2^d$, $\eta = \ f\ _{U^d}$, $\rho = \ f\ _{U^2}$ (<i>a norm</i>), $\sigma^2 = \mathbb{E}f^2$, $\zeta = \max f $; 14: $Q_3(W) = \ W\ _{U^3}^8$, $u = \ f\ _{U^3}$, $S = \ f\ _{U^2}^4$, $M = \sum a_r ^8$, $\rho = a_\xi $ and $b = a_{2\xi} $ (<i>half of 04's a and b</i>), D a tail fourth mass, $\eta = S_0 - S$, $t = u/\delta$, β_* (= half of 04's c), $C_* = (9/4)^{1/3}$, $\kappa = 2^{5/4}/12$. All are 03's and 04's Q_d .
moments	13: T_5 (= 04's B_5 for real f), B_1, B_2, B_3 six-point orbit moments (<i>not 04's or 06's B_k</i>), T_7 , $c_{d,p}(j)$ coloopleless counts (= 08's survivor coefficients), P_d (= 03's R_d). 14: $T(s, t)$ (= 04's), B_5 (= 04's), $M_\delta(u)$ <i>over all groups of order coprime to 6 and norm at most u</i> (04's (Q1) uses primes above 16 and norm exactly u ; Remark W4).
intervals, progressions	13: $D_N(M)$ the exact interval constant (<i>not a difference set</i>), $b = \min(\beta, 1 - \beta)$, $S_k(\delta) = \sum_{j \leq k-3} \delta^j$ (= 03's S_{k-2}).
colourings	14: $k = 2m$, D_ℓ , $s_\ell = \lceil D_\ell/2 \rceil$, $B = \sum s_\ell^2$ (= 10's B), q a field size (<i>10 works over primes</i>), $S_z(Q)$ symmetric directions, A_k^{sym} ordered counts.

Added in the fourth write (6 October 2026): sources 15–20. Contents. Section III.39: source 16, delivered Section 8 (fourth and fifth cube coefficients in every dimension; sharp quintic-moment bounds; optimality of the exponent $4/2^d$ for sets). Section III.40: source 18, delivered Section 3 (the expansion through order five on every finite abelian group, small minors and Smith forms, the centered cube bound, sharpness for sets). Sections III.41–III.48 are source 20's chapter: a guide and its delivered Sections 2–8 (directional cubes; two marked functions; the sharp directional-to-global comparison; second-order telescoping and quadratic counting; pair correlations; supersaturation; the sharp three-term obstruction).

Overlaps (dated notes). The expansion through order five: source 09 (on every group), sources 16 and 18 again; the quartic part also 03 and 08. Sharpness for sets: 03, 04, 09, 16, 18. Two-function counting and quadratic error: source 06's Section III.5 first, source 20 general and sharper; torsion: source 10, source 20 sharp per subgroup tuple. **Cross-source answers.** Source 06's Question 7 is answered for the norm-conversion step and its Question 6 advanced by source 20 (items III.9, III.8); source 18's question on odd-order coefficients is settled at $d = 3$ by source 04 (item III.1).

Reading conventions for the sections of sources 15–20 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
cube coefficients	16: $m = 2^d$, $\mathcal{Q}_d(\delta + tf)$ (= 03's and 04's Q_d), P_d (= 03's R_d), T_d , F_d , H_d (= 09's Q_d , S_d), $Q(f) = \ f\ _{U^2}^4$, Q_2 its order-two part, D (= 04's B_5 , 09's M_5), D_3 (= 09's $T_{3,5}$). 18: ℓ or d the cube dimension, $\mathcal{C}_\ell(g) = \ g\ _{U^\ell}^{2^\ell}$ (= 03's Q_d), R_ℓ , B_ℓ , Q_ℓ , S_ℓ (<i>18's Q_ℓ counts circuits; it is not a cube count</i>), V_4 , E_2 , T_5 , E_3 .
counting	20: Λ_k , Λ_L progression and system averages, “admissible” (multiplication by $1, \dots, k-1$ invertible), $D_{H_1, \dots, H_s}$ directional seminorms, κ_i torsion costs, $\sigma = \sqrt{\delta(1-\delta)}$, $R_m(\delta)$, $B_k(\delta)$ (<i>not 06's B_k</i>), $u = \ 1_A - \delta\ _{U^{k-1}}$, w_{ij} , $Q_d(f, g)$ a pair correlation, T_k^* (as 03's).

Added in the fifth write (6 October 2026): sources 21–27. Contents. Sections III.49–III.59 are source 21's chapter: a guide, its delivered Sections 2–9 (Fourier identities and the quartic deficit; the sixth- and seventh-degree terms with a retained tail; the order-five resonance; the sharp optimization and phase envelope; rigidity; the next coefficient on $\mathbb{Z}/5\mathbb{Z}$; indicator sharpness; a projection consequence of cube Cauchy–Schwarz), then source 23's delivered Section 7 and Appendix A (the same theorem again, a five-point model). Sections III.60–III.65 are source 22's chapter: a guide and its delivered Sections 6–10 (polynomial zero sets and prescribed equality patterns; progression specializations; quadratic classification; why additive blocks cannot extend the norm colouring; a global obstruction in characteristic-divisible degrees).

Overlaps (dated notes). The exponent-five coefficient: source 14's Theorem III.35.2 (exponent six and the lower model), then sources 21 and 23 independently (the sharp coefficient). The quartic gap is source 04's

Theorem III.3.4. Polynomial colourings: source 10's Section III.30, source 14's Section III.36, then source 22 (every degree, every pattern; Warning's second theorem and the Delsarte–Goethals–MacWilliams theorem are classical and credited). **Cross-source answers.** Source 14's Question 1 is answered by sources 21 and 23 (item III.23); its Question 7 is advanced and its Question 8 has a first instance by source 22 (items III.19, III.20).

Reading conventions for the sections of sources 21–27 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
exponent five	21: $Q_3(W) = \ W\ _{U^3}^8$, a_r Fourier coefficients, $\lambda_r = a_r ^2$, $S = \ f\ _{U^2}^4$, $M = \sum a_r ^8$, $T(s, t)$, $\mathcal{E}(\lambda)$, ρ , D the fourth-power tail (<i>not a dimension</i>), B_5 (= 04's), C_6 , C_7 (<i>not 01's or 06's C_k</i>), H , J , K six-vertex sums, $b = a_{2\xi} $, $\theta = 2\alpha - \beta$, $\phi = 3\alpha + \beta$, C_5 , β_5 , $\eta = S_0 - S$, $t = u/\delta$, $M_{\delta,5}(u)$. 23: $Q(W)$, A_6 , B_6 , C_6 (<i>23's C_6 is one of three forms; 21's C_6 is their sum</i>), $\kappa = 2^{1/4}/6$ (= 21's β_5), $\theta = \arg a_\xi$ (<i>not 21's θ</i>).
colourings	22: q the field size, $Q: \mathbb{F}_q^n \rightarrow \mathbb{F}_q^r$ a colouring with degree caps D_a , Π an equality pattern, $S_z(Q)$ successful directions, $B(\Pi, D)$, B_{sym} , B_{mono} budgets (= 14's B ; <i>not 18's B_d</i>), N_t norm forms, $T_b(R)$ third-difference supports, E_t field extensions.

Added in the sixth write (6 October 2026): sources 28–33. Contents. Sections III.66–III.70: a guide, source 29's delivered Sections 2–4 (Fourier identities; positive autocorrelation moments; consequences for the higher norm constants) and source 28's Section 8 (a stronger real norm gap and a higher-order obstruction).

Overlaps (dated notes). Source 28's universal bound is superseded by source 29's; the number 222 is computed by both, independently. **Cross-source answers.** Source 03's Research question 1 and source 04's Question 3 are advanced for $d \geq 4$ (item III.1); the sub-question whether real parts of characters are extremal is answered (no) for $d \geq 4$ by source 28.

Reading conventions for the sections of sources 28–33 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
odd constants	29: $c_f(h)$ the autocorrelation, $\lambda_\chi = \widehat{f}(\chi) ^2$, $S = \ f\ _{U^2}^4$, $m = 2^{d-2}$, C_m , B_d , $K_m(\tau)$, τ the self-inverse fraction, P_d the quartic cube coefficient (= 03's R_d on odd groups). 28: the same c_d^{odd} , f_b , F two-harmonic functions, $R_d(b)$ a ratio (<i>not 03's R_d</i>), $P(b)$ a polynomial, p the largest-pair mass (<i>not a prime</i>).

Added in the seventh write (6 October 2026): sources 34–39. Contents. Sections III.71–III.72: source 34's Section 10 (a stronger fourth-dimensional norm comparison) and source 35's Section 6 (a stronger gap between real centered uniformity norms), continuing the chapter on c_d^{odd} . Sections III.73–III.81: a guide and source 34's Sections 2–9 (Fourier constraints and the row quotient; the determinant reduction; exact first cases; a positive detector; consecutive binary determinants; optimal cyclic support; universal positivity; the exact quartic term). Sections III.82–III.86: a guide and source 38's Sections 2–5 (Fourier identities; exact reduction to a five-point quotient; the five-point profile; cyclic reduction for every fixed odd exponent). Sections III.87–III.88: a guide and source 38's Section 7 (an exact six-term extension defect).

Overlaps (dated notes). Source 34's lower example, dual identity and single-pair non-extremality are source 28's; sources 34 and 35 are independent same-day proofs of better coefficients, source 34's the stronger. **Cross-source answers.** Source 09's prime-onset question is answered by source 34 (item III.16); three of source 21's questions by source 38 (item III.27); item III.1 is advanced to $0.5587725 \leq c_4^{\text{odd}} \leq 2/3$.

Reading conventions for the sections of sources 34–39 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
odd constants	34: $a_\xi = \widehat{f}(\xi)$, $S = \ f\ _{U^2}^4$, p_j pair masses and p the largest (<i>not a prime</i>), C_m (= 29's), $H_m(p)$, D_4 the fourth dual operator, ϕ_j real characters, $K = 222^{1/4}/2$, $L = 27q_3^{1/4} + q_5^{1/4}$, β_d coupled bounds. 35: $\lambda_r = a_r ^2$, S , $M = \sum \lambda^4$ (<i>not 35's M of Part II</i>), $J(\lambda)$, A , B , $D = 1 - p$ (<i>not a divisor</i>).
prime torsion	34: M_S the augmented Boolean matrix, A_S its row quotient, $h(p)$, $h_{\text{cyc}}(m)$, $h_{\text{pos}}(m)$, $D_{\text{div}}(p)$, $\text{row}_{\mathbb{Z}}$; $h(p) = 09$'s $h(\ell)$.
exponent five	38: $\mathcal{C}_3(W) = \ W\ _{U^3}^8$ (<i>not $\mathbb{Q}$</i>), $\mathcal{M}_{\delta,G}(u)$, $u = \ f\ _{U^3}$, c_{10} , c_{12} , m an odd exponent; its $\mathcal{M}_{\delta,G}$ is 21's $M_{\delta,G}$.
norm colourings	38: N_d norm maps, e_2 the quadratic coefficient of the quintic norm, A_6^{sym} the six-term count, R the independent addition (<i>not a radius</i>).

Added in the eighth write (7 October 2026): sources 40–46. Contents. Sections III.89–III.91: a guide and source 45's Sections 4–5 (a quantitative real fourth-uniformity gap; applications to cube counting and structured holes), continuing the chapter on c_d^{odd} .

Overlaps (dated notes). Source 45's single-pair non-extremality and dual identity are sources 28's and 34's, its moment bound sources 29's and 34's, its three-pairing count that of sources 21, 35 and 38; it cites none of them. **Cross-source answers.** Item III.1 is advanced to $0.5587725 \leq c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$.

Reading conventions for the sections of sources 40–46 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
odd constants	45: $Q_d = \ f\ _{U^d}^{2^d}$, $S = Q_2$ (= 29's S), v , w the mean and the rest of $Q_2(g)$, $t(h)$, $w(h)$ derivative statistics, V , W their averages, $P(t)$ the scalar certificate, $\kappa = (832/4795)^{1/4}$ (<i>not 34's or 38's κ</i>), ϕ_j real characters, $\mathcal{D}_4$ the dual operator (= 34's D_4), $P_d = (6^d - 2 \cdot 4^d + 2^d)/8$ (= 04's quartic coefficient).

Added in the ninth write (7 October 2026): sources 47–53. Contents. Sections III.92–III.93: a guide and source 49's Section 7 (two-harmonic dual tests and the real uniformity gaps), continuing the chapter on c_d^{odd} and source 45's chapter.

Overlaps (dated notes). Source 49's balanced-moment profile is source 34's (it says so) and its test function grows out of source 34's two-harmonic example; it does not know source 45. **Cross-source answers.** Item III.1 is advanced to $0.5587725 \leq c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$, and $c_5^{\text{odd}} \leq 29/50$.

Reading conventions for the sections of sources 47–53 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
odd constants	49: p the largest inverse-pair mass (<i>not a prime</i>), $S = \ f\ _{U^2}^4$ (= 29's S), C_m , $H_m(p)$ (= 34's), ϕ_k real characters, $g_t = \phi_1 + t\phi_3$ the test, D_4 its dual (= 34's), $d_{k,n}(t)$ dual coefficients, M , E , W certificate constants (<i>not 45's M_m or W</i>), η a tail ratio.

Added in the eleventh write (7 October 2026): sources 60–69. Contents. Sections III.94–III.99: a guide, source 63's Sections 2–5 (conventions; the sharp bounded inequality; the high-order odd-group constant; the quartic cube coefficient) and source 60's Section 2 (the limiting real Gowers norm constant). Sections III.100–III.107: a guide and source 64's Sections 2–8 (weighted separation; no perfect cyclic alphabet; the nine-copy transfer; binomial transfer to the torus; endpoint exponents; deterministic realization; the dimension barrier). Sections III.108–III.109: a guide and source 67's Section 6 (a degree-free obstruction for

independent norm layers). **Cross-source answers.** The limit question of sources 28 and 29 and source 60's Research questions 1–3 (first part) are answered by source 63 (item III.1).

Reading conventions for the sections of sources 60–69 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Norm constants	63, 60: c_d^{odd} as in Section III.66 (the ratio of fourth powers); 63: $b(G)$, $c_d(G)$ on a fixed group, $P_d = (6^d - 2 \cdot 4^d + 2^d)/8$ (= 04's P_d), A_d the optimal quartic coefficient (<i>not 55's cube count A_d</i>), $\mathbb{R}/\mathbb{Z} = \mathbb{R}/\mathbb{Z}$; 60: β_{odd} , $\beta(q)$, $c_d(q)$, $C(d, P)$, $p_*(G)$ the least prime divisor, $B = 1/2 - 8/\pi^2 + 64/\pi^4$.
Section 4	64: Λ_k the progression count, $\rho_k(\alpha)$ (Deng–Tidor–Zhao's), β_6 the growth exponent, $b_6 = \log_4 80$, $v_6 = 5 + 9/b_6$, D an alphabet, w weights, B a modulus, V_k , W_k constants (<i>not 63's B or 69's B_d</i>); 67: $N_{\mathbb{F}_{q^d}/\mathbb{F}_q}$ norms, $M_{d,m}(q)$ centre counts, A_{2m}^{sym} progression counts, r output coordinates (<i>not 22's or 10's r in the same role only up to notation</i>).

III.1 Progression counting with fewer error terms

Source 03, delivered Section 5.

For integers $a_1, \dots, a_k$ and functions $g_i: G \rightarrow \mathbb{C}$, define

$$\Lambda_a(g_1, \dots, g_k) = \mathbb{E}_{x, t \in G} \prod_{i=1}^k g_i(x + a_i t).$$

We say that the slopes have invertible differences on G if each map $x \mapsto (a_i - a_j)x$, $i \neq j$, is an automorphism. For $G = \mathbb{Z}/N\mathbb{Z}$ this means $\gcd(a_i - a_j, N) = 1$. Consecutive slopes satisfy the condition whenever $\gcd(N, (k-1)!) = 1$, which is weaker than primality of N .

Lemma III.1.1 (Generalized von Neumann inequality). *Suppose the slopes have invertible differences and $k \geq 2$. For $|F_i| \leq 1$,*

$$|\Lambda_a(F_1, \dots, F_k)| \leq \|F_k\|_{U^{k-1}(G)}.$$

The same bound holds with any one of the functions distinguished.

Proof. For $k = 2$, the map $(x, t) \mapsto (x + a_1 t, x + a_2 t)$ is a bijection, so the average factors and is at most $|\mathbb{E} F_2| = \|F_2\|_{U^1}$. For the inductive step replace x by $y - a_1 t$, apply Cauchy–Schwarz in y , and remove the bounded factor $F_1(y)$. On expanding the square, write the two step variables as t and $t + u$. With $D_h F(z) = F(z) \overline{F(z+h)}$, the resulting expression is

$$\mathbb{E}_u \mathbb{E}_{y, t} \prod_{i=2}^k D_{(a_i - a_1)u} F_i(y + (a_i - a_1)t).$$

Its absolute value is at most $\mathbb{E}_u \|D_{(a_k - a_1)u} F_k\|_{U^{k-2}}$ by induction. Hölder and the invertibility of $a_k - a_1$ bound this by

$$\left(\mathbb{E}_h \|D_h F_k\|_{U^{k-2}}^{2^{k-2}} \right)^{1/2^{k-2}} = \|F_k\|_{U^{k-1}}^2.$$

This bounds the square of the original average. Permuting the indices allows any function to be distinguished. $\square$

Theorem III.1.2 (Density-weighted telescoping). *Assume the slopes have invertible differences. Let $0 \leq g_i \leq 1$, $\delta_i = \mathbb{E}g_i$, and $f_i = g_i - \delta_i$. For any permutation π of $\{1, \dots, k\}$,*

$$\left| \Lambda_{\mathbf{a}}(g_1, \dots, g_k) - \prod_i \delta_i \right| \leq \sum_{j=3}^k \left(\prod_{\ell=j+1}^k \delta_{\pi(\ell)} \right) \|f_{\pi(j)}\|_{U^{j-1}}. \quad (\text{III.1.1})$$

For $k = 2$ the error is zero.

Proof. Expand by replacing the functions with their means one at a time, starting with the last function in the chosen order. The exact difference is a sum of k terms. Term j has coefficient $\prod_{\ell>j} \delta_{\pi(\ell)}$ and an average containing $g_{\pi(1)}, \dots, g_{\pi(j-1)}, f_{\pi(j)}$. The first term vanishes because $f_{\pi(1)}$ has mean zero. The second vanishes by the two-variable bijection used in Lemma III.1.1. Apply that lemma to each remaining term, with its j slopes. $\square$

[write] **6 October 2026; sources 11–14.** Source 13's Theorem III.33.2 and Corollary III.33.3 prove this telescope and criterion again (its S_k is S_{k-2} here), with an explicit $\mathbb{Z}/3\mathbb{Z}$ example showing that the invertibility hypothesis of the generalized von Neumann bound cannot simply be dropped.

In the setting of Corollary 3.3, use the original index order. The normalized uniformity norms that occur in (III.1.1) are then at most its parameter α . Therefore its bound $2^k \alpha$ may be replaced by

$$\alpha \sum_{j=3}^k \prod_{i=j+1}^k \delta_i \leq (k-2)\alpha.$$

The operation is a telescoping expansion, which has one discrepancy per replacement. It avoids the unnecessary 2^k -term expansion of all balanced factors simultaneously. This is a standard multilinear technique; its role here is to record the stronger bound with the exact inhomogeneous parameters used by the source.

III.1.1 Optimizing the order of replacement

If all available norms in (III.1.1) are bounded by the same η , arrange the densities in nonincreasing order. This minimizes the displayed upper bound: every tail product then consists of the smallest densities of the required size. When distinct functions have different norms at different levels, define for a set of indices S

$$B(S) = 0 \quad (|S| \leq 2), \quad B(S) = \min_{i \in S} \{ \|f_i\|_{U^{|S|-1}} + \delta_i B(S \setminus \{i\}) \} \quad (|S| \geq 3).$$

This finite recursion computes the best bound among these telescoping orders in $O(k2^k)$ arithmetic operations once the norms are supplied. It is an optimization of the bound, not a claim that the bound is the exact counting error. For fixed small k , this separates analytic estimates from a finite, reproducible choice of order.

Corollary III.1.3 (An explicit modular progression criterion). *Let $k \geq 3$, let $|G| = N$, and suppose multiplication by each of $1, \dots, k-1$ is an automorphism of G . Let $A \subset G$ have density $0 < \delta \leq 1$, and set*

$$\eta = \|1_A - \delta\|_{U^{k-1}}, \quad S_{k-2}(\delta) = 1 + \delta + \dots + \delta^{k-3}.$$

The number of ordered pairs (x, t) with $t \neq 0$ and $x, x+t, \dots, x+(k-1)t \in A$ is at least

$$N^2(\delta^k - \eta S_{k-2}(\delta)) - \delta N. \quad (\text{III.1.2})$$

In particular, if

$$\eta \leq \frac{\delta^k}{2S_{k-2}(\delta)}, \quad N \geq 4\delta^{1-k}, \quad (\text{III.1.3})$$

there are at least $\delta^k N^2/4$ such pairs. Each has k distinct vertices. In the notation of the source, it suffices that

$$\alpha \leq \left(\frac{\delta^k}{2S_{k-2}(\delta)} \right)^{2^{k-1}}$$

when A is α -uniform of degree $k - 2$.

Proof. Gowers norm monotonicity gives $\|1_A - \delta\|_{U^{j-1}} \leq \eta$ for $3 \leq j \leq k$. Apply Theorem III.1.2 to equal functions and consecutive slopes. Exactly $|A| = \delta N$ parameter pairs have $t = 0$. For $t \neq 0$, equality of vertices with indices i, j would imply $(i - j)t = 0$, contrary to the automorphism assumption. Under (III.1.3), the right side of (III.1.2) is at least $\delta^k N^2/4$. $\square$

Remark III.1.4 (Modular and integer progressions). Corollary III.1.3 concerns progressions in the group. They do not automatically become ordinary integer progressions under the representative map to $\{0, \dots, N - 1\}$. One sufficient rectification condition is that every representative vertex lies in an interval of length H with $2(H - 1) < N$. Its consecutive second differences are zero modulo N and have absolute value below N , so they are zero as integers. This proves rectification, but obtaining the necessary interval control is additional information. The general modular counting criterion by itself supplies no such localization.

[write] **6 October 2026.** This rectification argument is formalized, in the form of source 04's Lemma III.4.5 (note there): Lean's `hasNatAP_of_short_modular_sequence` takes representatives in $\{0, \dots, n\}$ with $2n < N$. The remark's interval of H consecutive integers reduces to that case by a translation, a step not stated in Lean. Sections III.1 and III.4 prove the same telescoping bound independently; see the note after Theorem III.4.3.

Remark III.1.5 (Why small torsion cannot simply be ignored). For $G = \mathbb{Z}/2\mathbb{Z}$, slopes 0, 1, 2 have a noninvertible difference. Taking the nontrivial real character f , the average $\mathbb{E}_{x,t} f(x)f(x + 2t)$ equals one, although $\mathbb{E}f = 0$. The two centered factors are evaluated at the same point. Thus the pair cancellation used above really can fail when the slope hypothesis is removed.

III.2 The exact fourth-order term in the cube count

Source 03, delivered Section 6.

Gowers' Lemma 3.10 bounds a cube count by applying the triangle inequality to the Gowers norm. The resulting estimate is convenient, but it discards the mean-zero condition on the balanced function after that condition has been used to define uniformity. Retaining it has a substantial local effect: the first three error terms vanish, and the fourth can be computed exactly. The computation also identifies precisely what changes when the ambient finite abelian group has a nontrivial quotient of exponent two.

All results in this section concern the same labelled cubes as [1, Lemmas 3.7–3.10]; coincident vertices are allowed. The refinements are elementary consequences of the cube expansion, character orthogonality, and Gowers–Cauchy–Schwarz. Their proofs are included in full. They are useful sharpenings of the displayed lemma; we do not claim a new inverse theorem or a change in the best known quantitative bounds for Szemerédi's theorem.

III.2.1 Normalization and the centered expansion

Let G be a finite abelian group, with all averages normalized by the cardinality of their domain. Write $V_d = \{0, 1\}^d$, $m = 2^d$, and

$$L_v(x, h) = x + \sum_{i=1}^d v_i h_i \quad (v \in V_d).$$

For a real function $F : G \rightarrow \mathbb{R}$, define

$$Q_d(F) = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{v \in V_d} F(L_v(x, h)). \quad (\text{III.2.1})$$

Thus $Q_d(F) = \|F\|_{U^d}^m$; the usual conjugations in the definition of U^d have no effect for real functions. In particular, $Q_d(F) \geq 0$. If $F = \mathbf{1}_A$, the number of labelled cubes in A is $|G|^{d+1} Q_d(F)$.

Put $\delta = \mathbb{E}F$ and $f = F - \delta$. For $S \subseteq V_d$, put

$$M_S(f) = \mathbb{E}_{x, h} \prod_{v \in S} f(L_v(x, h)), \quad M_\emptyset(f) = 1. \quad (\text{III.2.2})$$

The exact finite expansion is

$$Q_d(\delta + f) = \sum_{S \subseteq V_d} \delta^{m-|S|} M_S(f). \quad (\text{III.2.3})$$

No limiting argument or regularity assumption is involved.

For real functions $g_v : G \rightarrow \mathbb{R}$, the normalized Gowers–Cauchy–Schwarz inequality gives

$$\left| \mathbb{E}_{x, h} \prod_{v \in V_d} g_v(L_v(x, h)) \right| \leq \prod_{v \in V_d} \|g_v\|_{U^d}. \quad (\text{III.2.4})$$

This is [1, Lemma 3.8] with normalized sums; its proof by successive Cauchy–Schwarz inequalities works on every finite abelian group. Taking $g_v = f$ on S and $g_v = 1$ elsewhere shows that, if $\eta = \|f\|_{U^d}$,

$$|M_S(f)| \leq \eta^{|S|}. \quad (\text{III.2.5})$$

We use also the standard monotonicity $\|f\|_{U^2} \leq \|f\|_{U^d}$ for $d \geq 2$.

Lemma III.2.1 (Cancellation through degree three). *If $\mathbb{E}f = 0$ and $1 \leq |S| \leq 3$, then $M_S(f) = 0$ on every finite abelian group.*

Proof. For $|S| = 1$, average over x . Otherwise choose a coordinate in which the distinct selected vertices are not all equal. For two or three vertices, one side of this coordinate contains exactly one selected vertex. Reflect that coordinate if necessary, so that this vertex is the only one with coordinate 1. A coordinate reflection preserves the average: it corresponds to the bijective change of variables $x' = x + h_i$, $h'_i = -h_i$, leaving the other h_j unchanged. After the reflection, h_i occurs in exactly one factor in the product defining M_S . Its average is $\mathbb{E}f = 0$. $\square$

Corollary III.2.2 (First refinement of Lemma 3.10). *For $d \geq 2$, $\delta \geq 0$, and real mean-zero f ,*

$$|Q_d(\delta + f) - \delta^m| \leq \sum_{j=4}^m \binom{m}{j} \delta^{m-j} \eta^j. \quad (\text{III.2.6})$$

Consequently the cube-count error is $O_{d, \delta}(\eta^4)$ as $\eta \rightarrow 0$ with d and positive δ fixed.

Proof. Expand using (III.2.3), discard the terms that vanish by Lemma III.2.1, and apply (III.2.5) to every remaining term. $\square$

The coefficient $\binom{m}{4}$ in (III.2.6) is itself far from necessary. Most four-vertex supports vanish as well.

III.2.2 Fourier energies and the quotient by doubling

Let $\widehat{G}$ be the character group, with characters written multiplicatively, and use normalized Fourier coefficients

$$\widehat{f}(\chi) = \mathbb{E}_{x \in G} f(x) \overline{\chi(x)}.$$

Write 1 for the trivial character. For real mean-zero f , set

$$\mathcal{E}(f) = \sum_{\chi \in \widehat{G}} |\widehat{f}(\chi)|^4 = \|f\|_{U^2}^4, \quad (\text{III.2.7})$$

$$\mathcal{E}_{(2)}(f) = \sum_{\substack{\chi \in \widehat{G} \setminus \{1\} \\ \chi^2 = 1}} |\widehat{f}(\chi)|^4. \quad (\text{III.2.8})$$

The omitted trivial-character term is zero. In particular,

$$0 \leq \mathcal{E}_{(2)}(f) \leq \mathcal{E}(f) \leq \eta^4. \quad (\text{III.2.9})$$

The second energy has a concrete quotient interpretation. Let $2G = \{2x : x \in G\}$ and

$$P_2 f(x) = \mathbb{E}_{z \in 2G} f(x + z).$$

This is the conditional average of f on the cosets of $2G$. A character survives this average exactly when it annihilates $2G$, equivalently when $\chi^2 = 1$. Fourier orthogonality therefore gives

$$\mathcal{E}_{(2)}(f) = \|P_2 f\|_{U^2(G)}^4. \quad (\text{III.2.10})$$

Thus the correction below measures the part of f visible in $G/2G$. If $|G|$ is odd, $2G = G$ and this correction is zero. If every element of G has order dividing two, then $P_2 f = f$.

III.2.3 A complete classification of four-vertex supports

Call a four-element subset $S \subseteq V_d$ an *ordinary parallelogram* if its vertices can be ordered a, b, c, e with $a + b = c + e$ as vectors in $\mathbb{Z}^d$. Call it a *binary tetrahedron* if

$$\bigoplus_{v \in S} v = 0$$

under coordinatewise addition modulo two, but S is not an ordinary parallelogram. This terminology describes the support in the Boolean cube, rather than a nondegeneracy condition on its image in G .

Lemma III.2.3 (The three possible fourth-order contributions). *For real mean-zero f and a four-element $S \subseteq V_d$,*

$$M_S(f) = \begin{cases} \mathcal{E}(f), & S \text{ is an ordinary parallelogram,} \\ \mathcal{E}_{(2)}(f), & S \text{ is a binary tetrahedron,} \\ 0, & \bigoplus_{v \in S} v \neq 0. \end{cases} \quad (\text{III.2.11})$$

These cases exhaust all four-element subsets.

Proof. If the XOR of the four vertices is nonzero, some coordinate has an odd number of ones, namely one or three. Reflect this coordinate when there are three. As in Lemma III.2.1, one difference variable then appears in exactly one centered factor, so the average vanishes.

Suppose that the XOR is zero. Reflect coordinates so that one chosen vertex becomes 0, and denote the other three vertices by u, v, w . The reflection preserves both the averages and the property of being an ordinary parallelogram. We have $u + v + w = 0$ over $\mathbf{F}_2$. Accordingly the only possible coordinate patterns across these three rows are

$$000, \quad 110, \quad 101, \quad 011.$$

At least two of the three nonzero patterns occur, because u, v, w are nonzero and pairwise distinct. Sum the independent variables h_i over coordinates of each of the three nonzero types, naming the sums r, s, t respectively. Every nonempty sum is uniform on G , and the nonempty sums are mutually independent.

If exactly two types occur, the four arguments, in an appropriate order, are

$$x, \quad x + r, \quad x + s, \quad x + r + s.$$

Their coefficient vectors form an ordinary parallelogram, and the average of the product of the four real f values is $\|f\|_{U^2}^4 = \mathcal{E}(f)$.

If all three types occur, the four arguments are

$$x, \quad x + r + s, \quad x + r + t, \quad x + s + t. \quad (\text{III.2.12})$$

There is no pairing with equal sums of coefficient vectors. For example, each possible pairing is ruled out by one of the three occupied coordinate types. Thus this is precisely the binary tetrahedron case.

It remains to evaluate (III.2.12). Substitute the Fourier expansion of f into all four factors, assigning characters $\chi_0, \chi_1, \chi_2, \chi_3$ in the displayed order. Averaging over r, s, t imposes

$$\chi_1\chi_2 = 1, \quad \chi_1\chi_3 = 1, \quad \chi_2\chi_3 = 1.$$

Hence $\chi_1 = \chi_2 = \chi_3 = \chi$ with $\chi^2 = 1$. The average over x then forces $\chi_0 = \chi$. The correlation is therefore

$$\sum_{\chi^2=1} \widehat{f}(\chi)^4.$$

For such characters $\chi(x) \in \{-1, 1\}$, so $\widehat{f}(\chi)$ is real. The trivial character has coefficient zero, and the last display equals (III.2.8). $\square$

Remark III.2.4 (Why integer rank alone is insufficient). The four vectors in (III.2.12), including the coefficient of x , form a matrix of determinant ± 2 . They are independent over $\mathbb{Q}$ but need not give independent uniform values on G . Discarding all rationally independent supports would incorrectly omit $\mathcal{E}_{(2)}(f)$. The three Smith types for a four-vertex support are $(1, 1, 1, 0)$, $(1, 1, 1, 2)$, and $(1, 1, 1, 1)$, corresponding respectively to an ordinary parallelogram, a binary tetrahedron, and the final case of Lemma III.2.3. In the first type the final zero records rank three; for $d = 2$ the rectangular matrix has only three nonzero invariant factors.

Lemma III.2.5 (Exact multiplicities). *The numbers of ordinary parallelograms and binary tetrahedra in V_d are respectively*

$$R_d = \frac{6^d - 2 \cdot 4^d + 2^d}{8}, \quad (\text{III.2.13})$$

$$T_d = \frac{8^d - 3 \cdot 6^d + 3 \cdot 4^d - 2^d}{24}. \quad (\text{III.2.14})$$

Their sum is

$$K_d = R_d + T_d = \frac{m(m-1)(m-2)}{24}. \quad (\text{III.2.15})$$

Proof. Count ordered quadruples satisfying $a + b = c + e$ in $\mathbb{Z}^d$. There are six possible coordinate patterns, so the total is 6^d . The degeneracies $a = c$, $b = e$ and $a = e$, $b = c$ each contribute 4^d , and their intersection contributes 2^d . These are all degeneracies: for example, $a = b$ forces $c = e = a$ coordinatewise. An ordinary parallelogram with four distinct vertices has exactly one pairing into two equal-sum pairs. Ordering both pairs and ordering the pairs gives eight ordered solutions. This proves (III.2.13).

For the XOR count, choose distinct ordered vertices a, b, c arbitrarily and put $e = a \oplus b \oplus c$. The fourth vertex differs from each of the first three. Every unordered four-element XOR-zero subset is counted 24 times. The total is thus $m(m-1)(m-2)/24$. Every ordinary parallelogram has XOR zero, and the other XOR-zero subsets are exactly the binary tetrahedra. Subtracting R_d gives (III.2.14). $\square$

For orientation, the first coefficients are

d	m	R_d	T_d	K_d
2	4	1	0	1
3	8	12	2	14
4	16	100	40	140
5	32	720	520	1240

For comparison, $\binom{m}{4}$ equals 70, 1820, 35960 when $d = 3, 4, 5$. Even after the first three terms have been removed, counting every fourth-order support by its absolute upper bound wastes most of the available cancellation.

III.2.4 The sharpened cube-count theorem

Theorem III.2.6 (Exact fourth-order expansion with explicit remainder). *Let $d \geq 2$, $m = 2^d$, let G be a finite abelian group, and let $F : G \rightarrow \mathbb{R}$ have mean $\delta \geq 0$. Put $f = F - \delta$ and $\eta = \|f\|_{U^d}$. Then*

$$Q_d(F) = \delta^m + \delta^{m-4}(R_d \mathcal{E}(f) + T_d \mathcal{E}_{(2)}(f)) + \mathcal{R}_d(F), \quad (\text{III.2.16})$$

where

$$|\mathcal{R}_d(F)| \leq \sum_{j=5}^m \binom{m}{j} \delta^{m-j} \eta^j. \quad (\text{III.2.17})$$

An empty sum is zero. In particular, the formula for $d = 2$ is exact with no remainder.

Proof. Use (III.2.3). The constant term is δ^m , and the terms of degrees one, two, and three vanish by Lemma III.2.1. Lemmas III.2.3 and III.2.5 give the total fourth-order term. The terms of degrees at least five satisfy (III.2.17) by (III.2.5) and the triangle inequality. $\square$

The formula holds in particular for arbitrary weights $0 \leq F \leq 1$, as well as for characteristic functions. The energies in the fourth term are nonnegative, even though the higher-order remainder can have either sign.

Corollary III.2.7 (Uniformity-parameter form). *Let $A \subseteq G$ have density δ , and suppose $\|\mathbf{1}_A - \delta\|_{U^d}^m \leq \alpha$. Then*

$$Q_d(\mathbf{1}_A) \leq \delta^m + K_d \delta^{m-4} \alpha^{4/m} + \sum_{j=5}^m \binom{m}{j} \delta^{m-j} \alpha^{j/m}. \quad (\text{III.2.18})$$

If $|G|$ is odd, the coefficient K_d in the fourth-order term can be replaced by R_d . The latter replacement also holds on an arbitrary group whenever $P_2(\mathbf{1}_A - \delta) = 0$.

Proof. Apply (III.2.9) and Theorem III.2.6. The vanishing of the quotient energy follows from (III.2.10). $\square$

In the notation of Gowers' Lemma 3.10, the old upper bound is $(\delta + \alpha^{1/m})^m$. For fixed d and positive δ , its displayed excess is of order $\alpha^{1/m}$; the new bound has excess of order $\alpha^{4/m}$. This comparison concerns the strength of the cube-count estimate under exactly the same norm hypothesis. It does not replace the arithmetic-progression counting or inverse arguments elsewhere in the proof.

For example, on an odd-order group and with $d = 3$,

$$Q_3(F) = \delta^8 + 12\delta^4 \|f\|_{U^2}^4 + \mathcal{R}_3(F), \quad (\text{III.2.19})$$

$$|\mathcal{R}_3(F)| \leq 56\delta^3 \eta^5 + 28\delta^2 \eta^6 + 8\delta \eta^7 + \eta^8. \quad (\text{III.2.20})$$

On an arbitrary group the exact fourth-order term has the additional $2\delta^4 \mathcal{E}_{(2)}(f)$.

If one prefers a single remainder expression, Taylor's theorem applied to $(\delta + t)^m$ gives, for $m \geq 5$,

$$\sum_{j=5}^m \binom{m}{j} \delta^{m-j} \eta^j \leq \binom{m}{5} \eta^5 (\delta + \eta)^{m-5}. \quad (\text{III.2.21})$$

The polynomial bound (III.2.17) usually retains more useful information when the density is small.

III.2.5 Sharpness and a necessary torsion correction

Proposition III.2.8 (The fourth power cannot be increased). *Fix $d \geq 2$ and $0 < \delta < 1$. Even if one restricts to odd-order finite abelian groups and weights $F : G \rightarrow [0, 1]$ of mean δ , there is no bound*

$$Q_d(F) - \delta^m = O_{d,\delta}(\|F - \delta\|_{U^d}^q) \quad (q > 4)$$

valid as the centered norm tends to zero. Among all finite abelian groups, K_d is the optimal leading coefficient in (III.2.18), expressed in terms of $\eta = \alpha^{1/m}$.

Proof. On $G = \mathbb{Z}/3\mathbb{Z}$, let $g(x) = \cos(2\pi x/3)$ and $F_t(x) = \delta + tg(x)$. For positive sufficiently small t , the weights lie in $[0, 1]$ and have mean δ . The two nonzero Fourier coefficients of g are $1/2$, so $\mathcal{E}(g) = 1/8$. Theorem III.2.6 gives

$$Q_d(F_t) - \delta^m = \frac{R_d}{8} \delta^{m-4} t^4 + O_{d,\delta}(t^5), \quad (\text{III.2.22})$$

where the remainder is identically zero when $d = 2$. Since $\|F_t - \delta\|_{U^d} = t \|g\|_{U^d}$ and $\|g\|_{U^d} > 0$, the exponent $q > 4$ is impossible.

For the coefficient over unrestricted groups, take a nontrivial character $\chi : G \rightarrow \{-1, 1\}$ and $F_t = \delta + t\chi$. Then $\|t\chi\|_{U^d} = t$ and $\mathcal{E}(t\chi) = \mathcal{E}_{(2)}(t\chi) = t^4$. Therefore

$$Q_d(F_t) - \delta^m = K_d \delta^{m-4} t^4 + O_{d,\delta}(t^5). \quad (\text{III.2.23})$$

Any smaller universal leading coefficient would fail for small t . $\square$

The same example proves that the torsion correction cannot be dropped. For a particularly transparent case, on $\mathbb{Z}/2\mathbb{Z}$ and with $d = 3$ one can compute the whole polynomial:

$$Q_3(\delta + t\chi) = \delta^8 + 14\delta^4 t^4 + t^8. \quad (\text{III.2.24})$$

Indeed, a subset S contributes precisely when both $|S|$ and every coordinate sum of its vertices are even. For the three-dimensional cube the only nonempty such supports have sizes four and eight: there are fourteen supports of size four and one of size eight. For size six, taking the complement would give two distinct vertices with equal XOR, which is impossible. Formula (III.2.24) directly distinguishes the correct coefficient 14 from the odd-order coefficient 12.

Remark III.2.9 (What sharpness has and has not shown). The identity determines the exact fourth-order coefficient in terms of the two Fourier energies. The constant K_d is also sharp when those energies are bounded only by η^4 over all groups. We have not asserted that R_d is the optimal leading constant in terms of η^4 when all groups are required to have odd order. That separate optimization involves the possible ratios $\|f\|_{U^2}^4 / \|f\|_{U^d}^4$ for real centered functions.

[write] 6 October 2026. *Settled for $d = 3$ by source 04.* By source 04's Theorem III.3.4, $\|f\|_{U^2}^4 \leq 2^{-1/2} \|f\|_{U^3}^4$ for real centered f on odd-order groups, with equality for $f = 2\operatorname{Re}(c\xi)$. Hence on odd-order groups the best leading constant in terms of $\eta^4 = \|f\|_{U^3}^4$ is $R_3 \cdot 2^{-1/2} = 6\sqrt{2} < 12 = R_3$ (source 04's Theorem III.3.6 and Proposition III.3.7); the proof is in the note after Theorem III.3.4. For $d \geq 4$ the question remains open (Research question V.14.1).

III.2.6 Sharpness can be transferred to actual sets

The use of small real weights in the previous proposition is not an essential restriction on the exponent. The following rounding argument records the quantitative reason.

Lemma III.2.10 (Rounding a bounded weight in a cube norm). *Let $p : G \rightarrow [0, 1]$ and $d \geq 2$, and choose a random subset A by including the elements of G independently with probabilities $p(x)$. For $e = \mathbf{1}_A - p$ and $m = 2^d$,*

$$\mathbb{E}_A \|e\|_{U^d}^m \leq \frac{\binom{m}{2}}{|G|}. \quad (\text{III.2.25})$$

Here $\mathbb{E}_A$ denotes expectation over the choice of the set, distinct from the averages over group elements.

Proof. Expand the Gowers norm as a cube average. When its m vertex values are pairwise distinct, independence and $\mathbb{E}_A e(x) = 0$ show that the expected product is zero. Every other product has absolute value at most one. For distinct vertices $u, v \in V_d$, the difference $L_u - L_v$ has a coefficient 1 or -1 on at least one of the independent variables h_i , and is therefore uniform on G . The probability of a collision between this pair is $|G|^{-1}$. The union bound over pairs proves (III.2.25). $\square$

Corollary III.2.11 (Optimal order and universal constant for sets). *For every fixed $d \geq 2$ and $0 < \delta < 1$, there are finite abelian groups G_n with $|G_n| \rightarrow \infty$ and subsets $A_n \subseteq G_n$ whose densities δ_n tend to δ , such that, writing $\eta_n = \|\mathbf{1}_{A_n} - \delta_n\|_{U^d}$, we have $\eta_n \rightarrow 0$ and*

$$\frac{Q_d(\mathbf{1}_{A_n}) - \delta_n^m}{\delta_n^{m-4} \eta_n^4} \rightarrow K_d. \quad (\text{III.2.26})$$

If the groups are required to have odd order, the power four is still optimal, although this argument does not identify the best leading constant in that restricted setting.

Proof. First take $G_n = (\mathbb{Z}/2\mathbb{Z})^n$, $N = |G_n|$, and a nontrivial real character χ_n . Put $t = N^{-1/(10m)}$ and $p = \delta + t\chi_n$, which lies in $[0, 1]$ for large n . Lemma III.2.10 and Markov's inequality give a realization of A_n with

$$\|\mathbf{1}_{A_n} - p\|_{U^d} \leq N^{-1/(2m)} = t^5$$

for all sufficiently large n , since the failure probability is at most $\binom{m}{2}N^{-1/2}$. The inequality $|\mathbb{E}e| \leq \|e\|_{U^d}$ implies $|\delta_n - \delta| \leq t^5$. After centering, the error in U^d is at most $2t^5$, so $\eta_n = t + O(t^5)$.

Both $\mathbf{1}_{A_n}$ and p have U^d norm at most one. The norm triangle inequality and $|a^m - b^m| \leq m|a - b|$ for $a, b \in [0, 1]$ show that

$$|Q_d(\mathbf{1}_{A_n}) - Q_d(p)| \leq mt^5.$$

The change from δ^m to δ_n^m is at most mt^5 as well. Combining these estimates with (III.2.23) proves (III.2.26).

For odd-order groups use $G_n = (\mathbb{Z}/3\mathbb{Z})^n$ and let g_n be the cosine of a nontrivial character. Its U^d norm is the same fixed positive number as the function g on $\mathbb{Z}/3\mathbb{Z}$ in Proposition III.2.8. Round $p = \delta + N^{-1/(10m)}g_n$ in exactly the same way. Equation (III.2.22) gives a strictly positive limit after division by η_n^4 . Consequently no exponent greater than four can hold uniformly even for characteristic functions on odd-order groups. $\square$

[write] **6 October 2026; sources 15–20.** Proved again by source 18 (Theorem III.40.12) and, for sets of exact density $1/3$ in odd cyclic groups, by source 16 (Proposition III.39.10).

Corollary III.2.12 (Exactly prescribed compatible rational densities). *If $0 < \delta < 1$ is rational, the sharp universal limit in (III.2.26) can be achieved with $\delta_n = \delta$ for every n . If the denominator of δ is odd, the odd-order sharpness of the fourth power can likewise be achieved at exactly that density.*

Proof. Write $\delta = a/q$ and use $G_n = (\mathbb{Z}/q\mathbb{Z}) \times (\mathbb{Z}/2\mathbb{Z})^n$, so $N = |G_n|$ makes δN an integer. Let χ_n be a nontrivial real character, put $t = N^{-1/(10m^2)}$, and round $p = \delta + t\chi_n$. The same Markov argument as before gives a realization A with

$$\|\mathbf{1}_A - p\|_{U^d} \leq N^{-1/(2m)} = t^{5m}.$$

Its density differs from δ by at most t^{5m} . Add or remove at most Nt^{5m} memberships to obtain a set B of cardinality exactly δN . If $e = \mathbf{1}_B - \mathbf{1}_A$ is supported on a fraction θ of the group, then $|e| \leq 1$ and

$$\|e\|_{U^d}^m \leq \mathbb{E}_{x,h} \left| \prod_v e(L_v(x, h)) \right| \leq \mathbb{E}_x \mathbf{1}_{\text{supp } e}(x) = \theta.$$

Thus the alteration costs at most t^5 in U^d norm. We obtain $\|\mathbf{1}_B - p\|_{U^d} = O(t^5)$, $\|\mathbf{1}_B - \delta\|_{U^d} = t + O(t^5)$, and $Q_d(\mathbf{1}_B) = Q_d(p) + O(t^5)$. Equation (III.2.23) gives the required limit. If q is odd, replace the second factor of G_n by $(\mathbb{Z}/3\mathbb{Z})^n$ and use the cosine of an order-three character. The same alteration argument and (III.2.22) prove optimality of the fourth power on odd-order groups at exactly the chosen density. $\square$

III.2.7 An exact formula for every higher coefficient

The fourth-order theorem is the first nontrivial instance of a complete Fourier description of the centered expansion. For a support $S \subseteq V_d$, define

$$\mathcal{K}_S(G) = \left\{ (\chi_v)_{v \in S} \in (\widehat{G} \setminus \{1\})^S : \prod_{v \in S} \chi_v = 1, \quad \prod_{v \in S} \chi_v^{v_i} = 1 \ (1 \leq i \leq d) \right\}.$$

All assigned characters are required to be nontrivial because f is centered. Character orthogonality gives the exact identity

$$M_S(f) = \sum_{(\chi_v) \in \mathcal{K}_S(G)} \prod_{v \in S} \widehat{f}(\chi_v). \quad (\text{III.2.27})$$

To prove it, expand each factor in (III.2.2); averaging over x gives the first character equation, and averaging over h_i gives the equation indexed by i . These equations are also sufficient for a Fourier term to survive. This is a finite exact formula, and it is valid without assumptions on the exponent or rank of G .

For groups of prime exponent, the equations can be studied as linear relations among the vectors $(1, v)$ over the prime field. For general finite abelian groups, their integer relation matrices and Smith invariant factors retain the additional torsion information. At four vertices, the only nontrivial invariant factor that can arise is two; this is why the single quotient energy in (III.2.8) suffices. Higher coefficients offer a concrete continuation: classify their support matrices, remove the supports for which (III.2.27) is empty, and estimate the remaining explicit Fourier correlations. The formula supplies an algorithmic route without assuming that the resulting classification will have a simple closed form.

III.2.8 The fifth-order remainder is also necessary

Proposition III.2.13 (Sharp order of the remainder). *Fix $d \geq 3$ and $0 < \delta < 1$. For weights $F : G \rightarrow [0, 1]$ on odd-order finite abelian groups, the remainder after the exact fourth-order term in Theorem III.2.6 cannot in general be bounded by $O_{d,\delta}(\|F - \delta\|_{U^d}^q)$ for any $q > 5$.*

Proof. Let $G = \mathbb{Z}/3\mathbb{Z}$, $g(x) = \cos(2\pi x/3)$, and $F_t = \delta + tg$ for positive sufficiently small t . Both nontrivial Fourier coefficients of g equal $1/2$, and its trivial coefficient is zero. Fourier expansion and character orthogonality show that every support moment $M_S(g)$ is nonnegative: each surviving term is a product of nonnegative Fourier coefficients.

In dimension at least three, consider the five-vertex support

$$S = \{0, e_1, e_2, e_3, e_1 + e_2 + e_3\}.$$

Its moment is

$$\mathbb{E}_{x,a,b,c} g(x)g(x+a)g(x+b)g(x+c)g(x+a+b+c) = \frac{1}{16}.$$

Indeed, assigning a nontrivial character χ to the last factor forces the three middle characters to be χ^{-1} and the first character to be χ^2 . Each of the two choices for χ therefore contributes $(1/2)^5$. Consequently

$$c_d := \sum_{|S|=5} M_S(g) \geq \frac{1}{16} > 0.$$

The exact finite expansion gives

$$\mathcal{R}_d(F_t) = c_d \delta^{m-5} t^5 + O_{d,\delta}(t^6),$$

whereas $\|F_t - \delta\|_{U^d} = t \|g\|_{U^d}$. Since $\|g\|_{U^d} > 0$, no exponent $q > 5$ can hold. $\square$

In dimension three the entire polynomial is especially simple:

$$Q_3(\delta + tg) = \delta^8 + \frac{3}{2}\delta^4 t^4 + \frac{1}{2}\delta^3 t^5 + \frac{1}{2}\delta^2 t^6 + \frac{1}{32}t^8. \quad (\text{III.2.28})$$

For a direct verification, put $a = \delta + t$ and $b = \delta - t/2$, the two possible weight values. Group the triples (h_1, h_2, h_3) by their number r of nonzero entries. Coordinate reflections allow each nonzero difference to be replaced by 1 when summing over x . For $r = 0, 1, 2, 3$ there are respectively 1, 6, 12, 8 difference triples, and their sums over x are respectively

$$a^8 + 2b^8, \quad 2a^4b^4 + b^8, \quad a^4b^4 + 2a^2b^6, \quad a^2b^6 + 2a^3b^5.$$

It follows that

$$81Q_3(F_t) = a^8 + 8b^8 + 24a^4b^4 + 32a^2b^6 + 16a^3b^5.$$

Substitution and expansion give (III.2.28). Its t^5 coefficient supplies a concrete witness to Proposition III.2.13.

[write] 6 October 2026. (III.2.28) is identical to source 04's (III.3.19), computed there by counting the 81 cube parameters by the number of vertices at 0; credited to both. The five-vertex support S above is the representative $\{000, 100, 010, 001, 111\}$ of source 04's eight surviving quintic supports (Lemma III.3.3), so $M_S(g) = B_5(g)$; the value $1/16$ is (III.3.8) with $a_{\pm 1} = 1/2$ and $2\xi = -\xi$ on $\mathbb{Z}/3\mathbb{Z}$.

Replacing t by $-t$ in (III.2.28) makes its fifth-order term negative. For sufficiently small positive t , this gives

$$Q_3(\delta - tg) < \delta^8 + 12\delta^4 \|tg\|_{U^2}^4.$$

Thus the exact quartic term is not, by itself, an unconditional lower bound for the whole excess; the sign of the remainder matters.

Remark W1 (The quartic lower bound with coefficient 12 is false; second write, 6 October 2026). For real $F = \delta + f$ with $\mathbb{E}f = 0$ on an odd-order group, the exact quartic term invites the lower bound

$$Q_3(F) \stackrel{?}{\geq} \delta^8 + 12\delta^4 \|f\|_{U^2}^4 \quad (0 \leq F \leq 1).$$

It is false. The display above shows this on $\mathbb{Z}/3\mathbb{Z}$ for sufficiently small t (source 03). Source 08 (Theorem III.10.1) and source 09 (Proposition III.22.2) prove it on $\mathbb{Z}/11\mathbb{Z}$ and on $\mathbb{Z}/q\mathbb{Z}$ for every prime $q \geq 17$, and source 08 gives an explicit instance with comfortable margins. *Counterexample.* On $G = \mathbb{Z}/11\mathbb{Z}$ let $h(x) = \cos(2\pi x/11) - \cos(4\pi x/11)$, $\delta = 1/2$, $\eta = 1/12$ and $F = \delta + \eta h$, so $1/3 \leq F \leq 2/3$. The nonzero Fourier coefficients of h are $\hat{h}(\pm 1) = 1/2$ and $\hat{h}(\pm 2) = -1/2$, so $\|\eta h\|_{U^2}^4 = \eta^4/4$, and the exact cube polynomial is

$$Q_3(\delta + \eta h) = \delta^8 + 3\delta^4\eta^4 - \frac{1}{2}\delta^3\eta^5 + \frac{3}{2}\delta^2\eta^6 + \frac{1}{4}\eta^8.$$

Hence $(Q_3(F) - \delta^8)/(\delta^4\|F - \delta\|_{U^2}^4) = 12 - 2s + 6s^2 + s^4$ with $s = \eta/\delta = 1/6$, which equals $15337/1296 \approx 11.834 < 12$. The polynomial is source 08's (III.10.3) and source 09's (III.22.4) at $\sigma = -1$. *Independent check at the write:* the polynomial was recomputed exactly by summing the Fourier products over all 11^4 free frequency choices of the eight cube vertices (exact rational arithmetic), and the value $Q_3(F) = 0.0039151672\dots$ was confirmed by direct floating-point summation over all 11^4 cubes. What survives is the lower bound $Q_3(F) - \delta^8 \geq 2\delta^4\|f\|_{U^2}^4$ (source 08's (III.12.2)); the best constant lies in $[2, 15337/1296]$ (source 08's question, Section V.2, item III.10).

III.2.9 A five-vertex correlation that detects order three

The next evaluated support illustrates why a higher-order classification will require more than the order-two energy. Consider the five vertices

$$S_3 = \{0000, 1110, 1101, 1011, 0111\} \subseteq \{0, 1\}^4.$$

For every real centered function f on a finite abelian group,

$$M_{S_3}(f) = \sum_{\substack{\chi \neq 1 \\ \chi^3 = 1}} \hat{f}(\chi)^4 \overline{\hat{f}(\chi)}. \quad (\text{III.2.29})$$

Here the four nonzero vertices are the all-ones vector with one entry removed. To prove (III.2.29), assign characters $\chi_0, \chi_1, \dots, \chi_4$ to the five factors, with χ_j corresponding to the vertex omitting coordinate j . Averaging over h_j imposes $\prod_{i \neq j, i \geq 1} \chi_i = 1$. Dividing any two such equations shows $\chi_1 = \dots = \chi_4 = \chi$; each equation then says $\chi^3 = 1$. The x -equation gives $\chi_0 = \chi^{-1}$. Fourier inversion and real-valuedness of f yield the formula.

This correlation vanishes if 3 does not divide $|G|$. On $\mathbb{Z}/3\mathbb{Z}$ with $f(x) = \cos(2\pi x/3)$ it equals $1/16$, and replacing f by $-f$ changes it to $-1/16$. In particular it is not a nonnegative energy of the fourth-order kind. We have evaluated this one support; we do not assert that it exhausts all possible fifth-order types.

[write] 6 October 2026; sources 06–10. *The fifth order is now classified (source 09).* Source 09's Theorem III.20.1 shows, on every finite abelian group and in every dimension, that a five-vertex centered moment can be nonzero only for two support types up to cube symmetry: the star $\{0, u, v, w, u + v + w\}$, whose moment is source 04's $B_5(f)$ ((III.3.5); source 08's S_5), and the ternary type $\{0, u + v + w, u + v + z, u + w + z, v + w + z\}$, whose moment is (III.2.29) (source 09's $T_{3,5}$). The support S_3 above is the ternary type with one coordinate per block; the support S in Proposition III.2.13 is a star. Source 09's Theorem III.21.1 gives the exact fifth-order coefficient $\delta^{m-5}(Q_d M_5 + S_d T_{3,5})$ with $Q_d = 2^d \left\{ \begin{smallmatrix} d+1 \\ 4 \end{smallmatrix} \right\}$ and $S_d = 2^d \left\{ \begin{smallmatrix} d+1 \\ 5 \end{smallmatrix} \right\}$. The four-vertex Smith types of Remark III.2.4 agree with source 09's Theorem III.19.1 (indices one and two only), and source 08's Lemma III.8.1 proves the same classification. At six vertices source 09's Theorem III.23.1 exhibits a support with row-lattice index five (checked at the write: $\det B = -5$, $M_{Sr} = (10, 5, 5, 5, 5, 5)$),

Smith form $(1, 1, 1, 1, 1, 5)$, so order-five torsion enters at order six. The exact quartic term of Theorem III.2.6 is also proved independently by source 08 (Theorem III.8.3) and source 09 (Theorem III.21.1), with the same coefficients $(R_d, T_d; \text{source 09 writes } P_d)$ and the same order-two energy.

III.3 Sharp centered refinements of the cube-counting lemma

Source 04, delivered Section 6.

Gowers' Lemma 3.10 bounds the number of d -dimensional cubes by applying the triangle inequality to the constant and balanced parts of an indicator. Centering contains additional information that this application of the triangle inequality discards. We retain it, first by identifying correlations that vanish, and then by proving a sharp comparison between the second and third uniformity norms of real centered functions. These results strengthen the stated cube-counting lemma; no historical priority beyond that comparison is asserted. As Gowers explains, this lemma is supplementary to his main density-increment argument, so an improvement here does not by itself imply an improved global bound in Szemerédi's theorem.

Throughout this section, averages on a finite abelian group G are normalized. For $d \geq 2$, put $q = 2^d$, $V_d = \{0, 1\}^d$, and

$$L_v(x, h) = x + v \cdot h, \quad Q_d(W) = \mathbb{E}_{x, h_1, \dots, h_d \in G} \prod_{v \in V_d} W(L_v(x, h)).$$

For real W , conjugation has no effect, so $Q_d(W) = \|W\|_{U^q}^q$. For an indicator, $|G|^{d+1}Q_d(1_A)$ is exactly Gowers' ordered cube count, including degenerate cubes. Write $W = \delta + f$, where $\delta \geq 0$, f is real, and $\mathbb{E}f = 0$; these are standing assumptions throughout the section. Set $u = \|f\|_{U^d}$. No bound on f is needed for the inequalities below.

III.3.1 Cancellation, parallelograms, and support counts

Lemma III.3.1 (Cancellation through degree three). *Any at most three distinct forms L_v are jointly uniform on the corresponding power of G , for every finite abelian group G . Consequently, if $\delta \geq 0$,*

$$|Q_d(\delta + f) - \delta^q| \leq \sum_{r=4}^q \binom{q}{r} \delta^{q-r} u^r. \quad (\text{III.3.1})$$

Proof. Reflect coordinates of the Boolean cube to move the first selected vertex to 0; these reflections correspond to invertible changes of the averaging variables. For three vertices the others are distinct nonzero Boolean vectors a, b . Two coordinates give a 2×2 minor of the matrix with rows a, b equal to 1 or -1 : use two coordinates distinguishing their supports, or one in their intersection and one in their difference if the supports are nested. Thus the map to $(x, x + a \cdot h, x + b \cdot h)$ is surjective, with equally sized fibers, on every abelian group. The cases of one or two vertices are immediate.

Expand the cube product according to the support on which f is chosen. Every term with one, two, or three factors of f vanishes. For a support of size r , the cube Cauchy–Schwarz inequality, with constant 1 at all other vertices, bounds its correlation in absolute value by u^r . Summation proves the assertion. $\square$

[write] **6 October 2026.** Lemma III.3.1 has the same statement and bound as source 03's Lemma III.2.1 with Corollary III.2.2 (this u is 03's η , this $q = 2^d$ is 03's m). Second route: source 03 uses coordinate reflections, this proof joint uniformity through 2×2 minors. Credited to both.

Odd group order permits a much more precise fourth-order term.

Proposition III.3.2 (Exact quartic coefficient). *Suppose $|G|$ is odd, and put $v = \|f\|_{U^2}$. Then*

$$Q_d(\delta + f) = \delta^q + P_d \delta^{q-4} v^4 + R_d, \quad (\text{III.3.2})$$

$$|R_d| \leq \sum_{r=5}^q \binom{q}{r} \delta^{q-r} u^r, \quad (\text{III.3.3})$$

$$P_d = 2^{d-3} (3^d - 2^{d+1} + 1). \quad (\text{III.3.4})$$

In particular, $P_2 = 1$, $P_3 = 12$, $P_4 = 100$, and $P_5 = 720$.

Proof. Four distinct Boolean vertices are either affinely independent or form a parallelogram. Indeed, any three are affinely independent. If four are dependent, their affine span is a plane, and some projection onto two coordinate axes is injective on that plane. The four images must be all four vertices of the unit square. The affine inverse of this projection preserves its parallelogram relation.

If the four vertices are independent, a 3×3 Boolean minor after this translation has determinant ± 1 or ± 2 . The determinant bound follows by adjoining a row and column of ones, changing binary entries b to $1 - 2b$, and applying Hadamard's inequality to the resulting 4×4 matrix. Multiplication by either determinant is invertible on an odd-order group. Thus the four forms are jointly uniform and their centered correlation is zero. On a parallelogram, any three forms are jointly uniform by Lemma III.3.1, while the fourth is their alternating sum. Its correlation is therefore exactly v^4 .

To count parallelograms, choose the k coordinates in which a pair of opposite vertices differs and fix the other coordinates. Around that center there are 2^{k-1} opposite pairs; choose two of them. The diagonal midpoint is unique, so

$$P_d = \sum_{k=2}^d \binom{d}{k} 2^{d-k} \binom{2^{k-1}}{2} = 2^{d-3} (3^d - 2^{d+1} + 1).$$

Bound the terms of degree at least five by cube Cauchy–Schwarz as before. □

[write] **6 October 2026.** Proposition III.3.2 is the odd-order case of source 03's Theorem III.2.6, which holds on every finite abelian group with the extra term $T_d \delta^{q-4} \mathcal{E}_{(2)}(f)$; that term vanishes on odd-order groups by (III.2.10). This P_d is source 03's R_d : $2^{d-3} (3^d - 2^{d+1} + 1) = (6^d - 2 \cdot 4^d + 2^d)/8$ (both give 1, 12, 100, 720). Second route: affine independence and Boolean minors $\pm 1, \pm 2$ in place of source 03's XOR classification. Do not read this remainder R_d as source 03's count R_d , or as the binomial sum $R_d(H)$ of Section IV.6 (Section F.3).

A general organizing principle is useful. Over $G = \mathbb{F}_p^n$, consider the vector matroid of the coefficient vectors $(1, v)$. If a selected support has a coloop, its centered correlation vanishes. Indeed, linear algebra gives a direction on which that form has coefficient 1 and all other selected forms have coefficient 0. Averaging along that direction averages one factor of f to zero. Thus only supports with no coloop need to be counted.

For $d = 3$ this count can be made completely explicit, and the result also holds on every odd-order abelian group.

Lemma III.3.3 (The surviving supports of a 3-cube). *Suppose that G has odd order. The numbers of supports with no coloop, in sizes $0, 1, \dots, 8$, are*

$$1, 0, 0, 0, 12, 8, 28, 8, 1.$$

All twelve quartic correlations equal $\|f\|_{U^2}^4$. The eight quintic correlations are equal to

$$B_5(f) = \mathbb{E}_{x,a,b,c} f(x)f(x+a)f(x+b)f(x+c)f(x+a+b+c). \quad (\text{III.3.5})$$

Consequently, with $u = \|f\|_{U^3}$ and $v = \|f\|_{U^2}$,

$$Q_3(\delta + f) = \delta^8 + 12\delta^4 v^4 + 8\delta^3 B_5(f) + R_6 + R_7 + u^8, \quad (\text{III.3.6})$$

$$|R_6| \leq 28\delta^2 u^6, \quad |R_7| \leq 8\delta u^7. \quad (\text{III.3.7})$$

Proof. The twelve coplanar quadruples are the six faces and six diagonal rectangles. A plane contains at most four Boolean vertices: some projection onto two coordinate axes is injective on the plane, and its Boolean image has at most four points. A five-set therefore has rank four and has a coloop exactly when it contains a coplanar quadruple. Each of the twelve quadruples has four extensions to a five-set. No five-set contains two such quadruples, since distinct planes meet in a line and a line contains at most two Boolean vertices. Thus $\binom{8}{5} - 12 \cdot 4 = 8$ five-sets survive. Every set of size at least six remains rank four after a deletion.

The eight surviving five-sets consist of one parity tetrahedron and one vertex of the opposite parity. Cube symmetries act transitively on them; a representative is $\{000, 100, 010, 001, 111\}$. The corresponding changes of averaging variables prove (III.3.5). These arguments remain valid on odd-order groups: any four independent cube forms have a minor ± 1 or ± 2 , hence are jointly uniform; an extra vertex outside a parallelogram is therefore independent of its three determining forms. Finally apply cube Cauchy–Schwarz to the 28 sextic and 8 septic supports. $\square$

For comparison, in characteristic two the support counts are $1, 0, 0, 0, 14, 0, 28, 8, 1$: the parity tetrahedra become two additional affine planes. The odd-order hypothesis is therefore a substantive restriction.

[write] **6 October 2026; sources 06–10.** *The complete three-cube polynomial in two more sources.* Source 08’s Theorem III.9.2 proves the exact identity behind (III.3.6) on every odd-order group, $Q_3(\delta + f) = \delta^8 + 12\delta^4 v^4 + 8\delta^3 S_5(f) + \delta^2(12A_6 + 12B_6 + 4C_6) + 8\delta A_7 + Q_3(f)$, with the same support counts $12, 8, 28, 8, 1$; its S_5 is B_5 here (the summands of the two Fourier forms are complex conjugates, and the sums are real). Source 09’s Theorem III.21.1 at $d = 3$ gives $\delta^4(12E_4 + 2T_2) + 8\delta^3 M_5$ on every finite abelian group; its T_2 vanishes in odd order. The characteristic-two counts 14 quartic supports above agree with source 08’s $F_3 = 14$ and source 09’s $A_3 = 14$. Sources 08 and 09 bound the quartic term by $12\delta^4 u^4$ in odd order; Theorem III.3.6 below has the sharp $6\sqrt{2}$.

III.3.2 A sharp gap between the real centered U^2 and U^3 norms

Use normalized Fourier coefficients

$$a_\xi = \widehat{f}(\xi) = \mathbb{E}_{x \in G} f(x) \overline{\xi(x)}, \quad \xi \in \widehat{G},$$

and write the dual group additively. Fourier inversion in (III.3.5) gives

$$B_5(f) = \sum_{\xi \in \widehat{G}} a_{2\xi} a_{-\xi}^3 a_\xi. \quad (\text{III.3.8})$$

For real centered f on an odd-order group, $a_{-\xi} = \overline{a_\xi}$, $a_0 = 0$, and the nonzero coefficients occur in distinct pairs. Therefore, with $\rho = \max_\xi |a_\xi|$,

$$|B_5(f)| \leq \rho \|f\|_{U^2}^4 \leq 2^{-1/4} \|f\|_{U^2}^5. \quad (\text{III.3.9})$$

Theorem III.3.4 (Sharp real centered norm gap). *Let G have odd order and let $f : G \rightarrow \mathbb{R}$ have mean zero. Then*

$$\|f\|_{U^3}^8 \geq 2 \|f\|_{U^2}^8. \quad (\text{III.3.10})$$

For nonzero f , equality holds exactly when

$$f(x) = 2 \operatorname{Re}(c\xi(x))$$

for a nontrivial character ξ and a nonzero complex number c .

Proof. We first use the Fourier identity

$$\begin{aligned} \|f\|_{U^3}^8 &= \sum_{s,t \in \widehat{G}} |T(s,t)|^2, \\ T(s,t) &= \sum_{r \in \widehat{G}} a_r \overline{a_{r+s}} \overline{a_{r+t}} a_{r+s+t}. \end{aligned} \quad (\text{III.3.11})$$

The Fourier coefficients are normalized, but all sums in (III.3.11) are unnormalized. To verify the identity, expand the original cube in Fourier series. The four orthogonality equations say precisely that its frequency variables form an affine cube:

$$r_v = r_{000} + v_1(r_{100} - r_{000}) + v_2(r_{010} - r_{000}) + v_3(r_{001} - r_{000}).$$

These equations are solved without division. Grouping the two parallel faces gives (III.3.11).

Put

$$\lambda_r = |a_r|^2, \quad S = \sum_r \lambda_r^2 = \|f\|_{U^2}^4, \quad M = \sum_r \lambda_r^4.$$

Retain only the two coordinate axes in the nonnegative sum (III.3.11). Since $T(0,t) = \sum_r \lambda_r \lambda_{r+t}$, this gives

$$\|f\|_{U^3}^8 \geq 2\mathcal{E}(\lambda) - S^2, \quad \mathcal{E}(\lambda) = \sum_{r+s=t+w} \lambda_r \lambda_s \lambda_t \lambda_w. \quad (\text{III.3.12})$$

The subtraction corrects the double counting of the origin.

The energy summands are nonnegative. Retain the union of three families:

$$\begin{aligned} \mathcal{A}: & \quad r = t, \quad s = w, \\ \mathcal{B}: & \quad r = w, \quad s = t, \\ \mathcal{C}: & \quad s = -r, \quad w = -t. \end{aligned}$$

Each family contributes S^2 , because λ is even. Each pairwise intersection contributes M . Their triple intersection has $2r = 0$; odd group order and $\lambda_0 = 0$ make its contribution zero. Thus

$$\mathcal{E}(\lambda) \geq 3S^2 - 3M,$$

and hence

$$\|f\|_{U^3}^8 \geq 5S^2 - 6M. \quad (\text{III.3.13})$$

Group nonzero frequencies into pairs $\{r, -r\}$, and let b_j denote λ_r^2 for a representative of each pair. Then

$$S = 2 \sum_j b_j, \quad M = 2 \sum_j b_j^2 \leq \frac{S^2}{2}.$$

This proves (III.3.10).

Equality in (III.3.10) forces equality in the last inequality, so only one Fourier pair can be nonzero. Conversely, a parallelogram on a two-point subset of an odd-order group is degenerate in at least one direction. If the Fourier support is $\{\xi, -\xi\}$, all nonzero terms in (III.3.11) therefore lie on its coordinate axes. Directly, the two norm powers are $8|a_\xi|^8$ and $4|a_\xi|^8$, respectively. $\square$

[write] 6 October 2026. *This answers source 03's Research question V.14.1 for $d = 3$ (cross-source finding of the intake; proof at the write).* Let G have odd order and $f : G \rightarrow \mathbb{R}$ be centered and nonzero; then $|G| \geq 3$, since on the trivial group a centered function vanishes. By (III.3.10),

$$\frac{\|f\|_{U^2}^4}{\|f\|_{U^3}^4} = \left(\frac{\|f\|_{U^2}^8}{\|f\|_{U^3}^8} \right)^{1/2} \leq 2^{-1/2},$$

with equality exactly for $f = 2 \operatorname{Re}(c\xi)$, ξ a nontrivial character and $c \neq 0$; such f exist on every odd-order group with $|G| \geq 3$. Hence $c_3^{\text{odd}} = 2^{-1/2}$, the supremum is attained, and the attaining functions are the real parts of characters. The best odd-order leading coefficient in source 03's reduction is $R_3 c_3^{\text{odd}} = 12 \cdot 2^{-1/2} = 6\sqrt{2}$, which is Theorem III.3.6 with Proposition III.3.7. This also settles source 03's Remark III.2.9 at $d = 3$: $R_3 = 12$ is not the optimal odd-order constant, $6\sqrt{2}$ is. For $d \geq 4$ the question is open; it is source 04's own Question 3 (Section V.2, item III.1).

[write] 6 October 2026; sources 11–14. Source 14 reproves this gap and the leading term $6\sqrt{2}$ (its Section III.34, equation (III.34.8)) on the way to the exact secondary coefficient.

[write] 6 October 2026; sources 15–20. Source 18 asks for the best odd-order fourth-order coefficient “especially in dimension four” (its Question 7) without citing this theorem, which settles $d = 3$. Source 16's Proposition III.39.8 reproves the bound on B_5 with $2^{-1/4}$ and adds $2^{-1/2}$ without three-torsion.

[write] 6 October 2026; sources 28–33. *Advanced for $d \geq 4$, not settled* (sources 28 and 29, Section III.66). Source 29 proves $\|f\|_{U^d}^{2^d} \geq 2^{-m} \binom{2m}{m} \|f\|_{U^2}^{2^d}$, $m = 2^{d-2}$, hence $c_4^{\text{odd}} \leq (8/35)^{1/4}$ and $c_d^{\text{odd}} \leq B_d \downarrow 1/2$, and $c_d^{\text{odd}} \geq 1/3$ for every d (Corollary III.69.1, Proposition III.69.2); source 28 proves $c_4^{\text{odd}} \geq 0.5587294$ and that a real part of a character is not extremal for $d \geq 4$ (Corollary III.70.8, Theorem III.70.9); the intake's certificate gives $c_4^{\text{odd}} \geq 0.5587725$. So $0.5587725 \leq c_4^{\text{odd}} \leq (8/35)^{1/4} = 0.6914416$ among sources 01–33; later placed sources record better upper bounds (item III.1).

[write] 6 October 2026; sources 34–39. *The upper end improved again* (sources 34 and 35, Sections III.71–III.72). Source 34 proves $\|f\|_{U^2}^4 \leq (2/3) \|f\|_{U^4}^4$, that is the coefficient $81/16$ in source 29's normalization (Theorem III.71.1); source 35, the same day, proves the coefficient 5 (Theorem III.72.1); for $5 \leq d \leq 8$ source 34's coupled bounds slightly improve source 29's (Remark III.71.10). So $0.5587725 \leq c_4^{\text{odd}} \leq 2/3$ among sources 01–39; source 45, placed and not yet written, records $(832/4795)^{1/4} = 0.6454070$ (item III.1).

[write] 7 October 2026; sources 40–46. *The upper end improved again* (source 45, Section III.89): $\|f\|_{U^4}^{16} \geq (4795/832) \|f\|_{U^2}^{16}$ (Theorem III.90.1), so $0.5587725 \leq c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$ among sources 01–46; source 49, placed by f8bc5e2ec and not yet written in, records $29/4$, that is $c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ (for its write) (item III.1).

[write] 7 October 2026; sources 47–53. *The upper end improved again* (source 49, Section III.92): $\|f\|_{U^4}^{16} \geq (29/4) \|f\|_{U^2}^{16}$ (Theorem III.93.1), so $0.5587725 \leq c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ among sources 01–53, and $c_5^{\text{odd}} \leq 29/50$ (item III.1).

Corollary III.3.5 (Stability of the norm gap). *Suppose that G has odd order. If $f \neq 0$, $0 \leq \varepsilon < 3$, and*

$$\|f\|_{U^3}^8 \leq (2 + \varepsilon)\|f\|_{U^2}^8,$$

then some nontrivial character ξ , with $c = \widehat{f}(\xi)$, satisfies

$$\|f - 2\operatorname{Re}(c\xi)\|_{U^2} \leq (\varepsilon/3)^{1/4}\|f\|_{U^2}. \quad (\text{III.3.14})$$

Proof. Equation (III.3.13) implies $M/S^2 \geq 1/2 - \varepsilon/6$. For the pair masses $p_j = 2b_j/S$, we have $\sum_j p_j = 1$ and $\sum_j p_j^2 \geq 1 - \varepsilon/3$. Since $\max_j p_j \geq \sum_j p_j^2$, one pair carries at least $1 - \varepsilon/3$ of the Fourier fourth-power mass. Removing it and using the Fourier formula for U^2 proves the claim. $\square$

This is a statement of stability in U^2 , not an assertion of approximation in L^2 . A large amount of Fourier square mass can be spread over many small coefficients.

III.3.3 The sharp leading cube bound and its models

Theorem III.3.6 (Refined form of Lemma 3.10 in dimension three). *Let G have odd order, $f : G \rightarrow \mathbb{R}$ have mean zero, $\delta \geq 0$, and $u = \|f\|_{U^3}$. Then*

$$\begin{aligned} Q_3(\delta + f) &\leq \delta^8 + 6\sqrt{2}\delta^4 u^4 + 2^{17/8}\delta^3 u^5 \\ &\quad + 28\delta^2 u^6 + 8\delta u^7 + u^8. \end{aligned} \quad (\text{III.3.15})$$

Consequently, if A is α -uniform of degree two in Gowers' convention,

$$\begin{aligned} Q_3(1_A) &\leq \delta^8 + 6\sqrt{2}\delta^4 \alpha^{1/2} + 2^{17/8}\delta^3 \alpha^{5/8} \\ &\quad + 28\delta^2 \alpha^{3/4} + 8\delta \alpha^{7/8} + \alpha. \end{aligned} \quad (\text{III.3.16})$$

Proof. In (III.3.6), use $v^4 \leq u^4/\sqrt{2}$ from Theorem III.3.4 and

$$|B_5(f)| \leq 2^{-1/4}v^5 \leq 2^{-7/8}u^5.$$

The remainder bounds give (III.3.15). Gowers' parameter satisfies $u^8 \leq \alpha$, which yields (III.3.16). $\square$

The leading error at fixed density is now $\alpha^{1/2}$, with an optimal constant, in place of the $\alpha^{1/8}$ term in $(\delta + \alpha^{1/8})^8$. In higher dimensions, monotonicity and Theorem III.3.4 similarly improve the quartic term in Proposition III.3.2 to $P_d \delta^{q-4} u^4/\sqrt{2}$ for $d \geq 3$; sharpness of that higher-dimensional constant is not established here.

Proposition III.3.7 (An exact model for the leading constant). *Let $p > 8$ be prime and $f(x) = t \cos(2\pi x/p)$ on $\mathbb{Z}/p\mathbb{Z}$. Then*

$$\|f\|_{U^2}^4 = t^4/8, \quad \|f\|_{U^3}^8 = t^8/32, \quad (\text{III.3.17})$$

$$Q_3(\delta + f) = \delta^8 + \frac{3}{2}\delta^4 t^4 + t^8/32 = \delta^8 + 6\sqrt{2}\delta^4 u^4 + u^8. \quad (\text{III.3.18})$$

For $0 < t \leq \min(\delta, 1 - \delta)$ this is a weight in $[0, 1]$.

Proof. The Fourier support is $\{1, -1\}$, with coefficients $t/2$. Odd-degree terms vanish, since an odd number of signs cannot sum to zero and $p > 8$ prevents modular aliasing. A six-set omits two distinct vertices. In a coordinate where those vertices differ, the selected set has exactly three ones, so no signed zero sum is possible there either.

In the full eight-vertex term, the positive-sign four-set must have two ones in each coordinate. There are exactly eight choices: six unions of two opposite-vertex pairs, and the two parity tetrahedra. To see completeness, if one opposite pair occurs, the remaining two vertices must also be opposite; otherwise, reflecting a selected vertex to 000 forces the other three to be 110, 101, 011. Thus the full term is $8(t/2)^8 = t^8/32$. The quartic term follows from Proposition III.3.2. $\square$

This proves that both the fourth power of u and its coefficient $6\sqrt{2}\delta^4$ are optimal. Moreover, a family attaining this leading coefficient as $u \rightarrow 0$ at fixed positive density must approach equality in Theorem III.3.4. Corollary III.3.5 then concentrates its Fourier fourth-power mass on one conjugate pair.

The next coefficient in (III.3.15) is also optimal over odd-order groups, but this assertion uses 3-torsion. On $\mathbb{Z}/3\mathbb{Z}$,

$$Q_3(\delta + t \cos(2\pi x/3)) = \delta^8 + \frac{3}{2}\delta^4 t^4 + \frac{1}{2}\delta^3 t^5 + \frac{1}{2}\delta^2 t^6 + t^8/32. \quad (\text{III.3.19})$$

For a direct verification, put $A = \delta + t$, $B = \delta - t/2$. Among the 81 cube parameters, the numbers with k vertices equal to zero are

k	0	2	3	4	8
number	8	32	16	24	1.

These counts follow by separating the cases of zero, one, two, or three nonzero increments. The cube count is therefore $(8B^8 + 32A^2B^6 + 16A^3B^5 + 24A^4B^4 + A^8)/81$; expansion gives (III.3.19). Since $u^8 = t^8/32$, its quintic term is exactly $2^{17/8}\delta^3 u^5$. This optimality is not asserted for large prime cyclic groups, where there is no 3-torsion.

[write] **6 October 2026.** (III.3.19) is identical to source 03's (III.2.28), computed there by grouping difference triples; credited to both.

III.3.4 Without 3-torsion: a sharp fractional correction exponent

The distinction just observed leads to a further refinement. The fifth-order correlation requires a frequency and its double. A function near equality in the sharp norm gap concentrates on a pair; without 3-torsion, doubling leaves that pair. Balancing the resulting loss and gain produces a fractional power.

Theorem III.3.8 (A correction of order $u^{16/3}$). *Suppose $|G|$ is coprime to 6, $\mathbb{E}f = 0$, f is real, and $\delta > 0$. With $u = \|f\|_{U^3}$,*

$$Q_3(\delta + f) \leq \delta^8 + 6\sqrt{2}\delta^4 u^4 + 6\delta^{8/3} u^{16/3} + 28\delta^2 u^6 + 8\delta u^7 + u^8. \quad (\text{III.3.20})$$

The exponent 16/3 cannot be increased after subtracting the sharp leading term. The coefficient 6 is an explicit upper bound, not a claimed optimal constant. In the parameter $\alpha = u^8$, the correction has order $\alpha^{2/3}$.

Proof. The zero function is immediate. Otherwise put $S = \|f\|_{U^2}^4$, choose a largest Fourier pair $\{\xi, -\xi\}$, and set

$$\rho = |a_\xi|, \quad z = 1 - \frac{2\rho^4}{S}, \quad D = zS = \sum_{r \notin \{\xi, -\xi\}} |a_r|^4.$$

Here $0 \leq z < 1$. Since $M = \sum_r |a_r|^8 \leq \rho^4 S = (1 - z)S^2/2$, (III.3.13) gives $u^8 \geq (2 + 3z)S^2$. The convex function $g(z) = (1 + 3z/2)^{-1/2}$ lies below its chord on $[0, 1]$, so

$$g(z) \leq 1 - (1 - \sqrt{2/5})z \leq 1 - z/3.$$

Consequently,

$$12\delta^4 S \leq 6\sqrt{2}\delta^4 u^4 - 2\sqrt{2}\delta^4 u^4 z. \quad (\text{III.3.21})$$

There is no 3-torsion, so 2ξ and -2ξ are outside the selected pair. The two selected terms of (III.3.8) contribute at most $2\rho^4(D/2)^{1/4}$ in absolute value; the other terms contribute at most ρD . Thus

$$\begin{aligned} |B_5(f)| &\leq 2^{-1/4} S^{5/4} ((1 - z)z^{1/4} + (1 - z)^{1/4} z) \\ &\leq 2^{3/4} S^{5/4} z^{1/4} \leq 2^{1/8} u^5 z^{1/4}. \end{aligned}$$

Combine this with (III.3.21). For $A, K > 0$,

$$\sup_{z \geq 0} (Kz^{1/4} - Az) = \frac{3K^{4/3}}{4^{4/3} A^{1/3}}.$$

Taking $A = 2\sqrt{2}\delta^4 u^4$ and $K = 2^{25/8}\delta^3 u^5$ makes the right side $6\delta^{8/3}u^{16/3}$. Inserting the remainder estimates from (III.3.7) proves (III.3.20). $\square$

Proposition III.3.9 (A two-frequency obstruction). *Fix $0 < \delta < 1$ and a prime $p > 16$. There are weights W_u on $\mathbb{Z}/p\mathbb{Z}$, with mean δ and $\|W_u - \delta\|_{U^3} = u$, such that, as $u \downarrow 0$,*

$$Q_3(W_u) - \delta^8 = 6\sqrt{2}\delta^4 u^4 + (9/4)^{1/3}\delta^{8/3}u^{16/3} + o_\delta(u^{16/3}). \quad (\text{III.3.22})$$

Proof. For $f(x) = a \cos(2\pi x/p) + b \cos(4\pi x/p)$, with $a, b > 0$, Fourier inversion gives

$$\|f\|_{U^3}^8 = \frac{a^8 + 6a^4 b^4 + b^8}{32}, \quad \|f\|_{U^2}^4 = \frac{a^4 + b^4}{8}, \quad B_5(f) = \frac{a^4 b}{16}. \quad (\text{III.3.23})$$

In fact there is the complete identity

$$\begin{aligned} Q_3(\delta + f) &= \delta^8 + \frac{3}{2}\delta^4(a^4 + b^4) + \frac{1}{2}\delta^3 a^4 b + \frac{3}{2}\delta^2 a^4 b^2 \\ &\quad + \frac{a^8 + 6a^4 b^4 + b^8}{32}. \end{aligned} \quad (\text{III.3.24})$$

Here is a finite counting proof of both this identity and the first formula in (III.3.23). Apply the Fourier cube identity (III.3.11) to $\delta + f$. Its frequency support is contained in $\{-2, -1, 0, 1, 2\}$, with coefficients $\delta, a/2, b/2$ at frequencies of absolute value 0, 1, 2, respectively. Since $p > 16$, the orthogonality equations have no modular aliasing, so the frequency variables are integer affine cubes in this five-point interval.

Reflecting directions with negative increments makes all increments nonnegative; their sum is at most four. Enumerate their nonzero magnitudes, choose their coordinate positions and signs, and

then choose the initial vertex between -2 and 2 minus their sum. The possibilities are the empty list, the singletons $(1), (2), (3), (4)$, the pairs $(1, 1), (1, 2), (1, 3), (2, 2)$, and the triples $(1, 1, 1), (1, 1, 2)$. If n_j is the number of vertices with frequency of absolute value j , this gives

(n_0, n_1, n_2)	number of frequency cubes
$(8, 0, 0)$	1
$(4, 4, 0)$	24
$(4, 0, 4)$	24
$(3, 4, 1)$	16
$(2, 4, 2)$	96
$(0, 8, 0)$	8
$(0, 4, 4)$	48
$(0, 0, 8)$	8

Each row contributes its count times $\delta^{n_0}(a/2)^{n_1}(b/2)^{n_2}$, proving (III.3.24). Setting $\delta = 0$ proves the stated U^3 formula. The remaining two formulas in (III.3.23) follow from the Fourier fourth moment and (III.3.8).

Let $c = (\sqrt{2}/6)^{1/3}$ and choose

$$b = c\delta^{-1/3}u^{4/3}, \quad a^4 = -3b^4 + \sqrt{32u^8 + 8b^8}.$$

For small u this is positive, $\delta + f$ takes values in $[0, 1]$, and (III.3.23) gives $\|f\|_{U^3} = u$ exactly. Since

$$a^4 = 4\sqrt{2}u^4 - 3b^4 + O(b^8/u^4),$$

equation (III.3.24) gives

$$Q_3(\delta + f) - \delta^8 = 6\sqrt{2}\delta^4u^4 - 3\delta^4b^4 + 2\sqrt{2}\delta^3u^4b + O_\delta(u^6).$$

Substitution of b produces the coefficient $-3c^4 + 2\sqrt{2}c = (9/4)^{1/3}$, proving (III.3.22). $\square$

For the local extremal problem over groups of order coprime to 6, these results determine the first coefficient and the next exponent. More precisely, subtract $6\sqrt{2}\delta^4u^4$ from the supremum of the cube excess over weights of mean δ with centered U^3 norm at most u , and divide by $\delta^{8/3}u^{16/3}$. The liminf of this normalized correction is at least $(9/4)^{1/3}$, and its limsup is at most 6. Proving convergence and determining the limit remain focused open questions; existence of an asymptotic second coefficient is not asserted here.

[write] 6 October 2026; sources 11–14. *Settled by source 14.* The limit exists and equals $(9/4)^{1/3}$, uniformly over groups of order coprime to 6 (Theorem III.34.1; Remark W4 for this normalization). The two-frequency family of Proposition III.3.9 is therefore asymptotically extremal; the 6 of Theorem III.3.8 stays an explicit, now superseded, upper bound.

III.3.5 Sharpness for indicator sets

Proposition III.3.10 (Rounding the sharpness examples). *For every fixed $0 < \delta < 1$, there are subsets A_j of prime cyclic groups, with densities $\delta_j \rightarrow \delta$ and $\alpha_j = \|1_{A_j} - \delta_j\|_{U^3}^8 \rightarrow 0$, such that*

$$Q_3(1_{A_j}) - \delta_j^8 \sim 6\sqrt{2}\delta^4\alpha_j^{1/2}. \quad (\text{III.3.25})$$

The obstruction to improving the exponent $16/3$ also persists for indicator sets on growing prime cyclic groups.

Proof. Independently include each $x \in \mathbb{Z}/p\mathbb{Z}$ with probability $w(x)$, where w is a prescribed weight of mean δ . At most $\binom{2^d}{2}/p$ of the normalized cube parameters identify two distinct vertices: each equality is one nontrivial linear equation with probability $1/p$. Away from these parameters, independence makes the expected product of indicators equal the product of their probabilities. The same argument applies to the centered variables $1_A(x) - \delta$. Thus the expectation error in the centered cube polynomial is at most $2\binom{2^d}{2}/p$.

Changing one inclusion decision changes a normalized cube polynomial by at most $2^d/p$. Indeed, telescope the product over its vertex occurrences; each occurrence has p^d possible remaining parameters, and the other factors are bounded by one. The bounded-difference variance inequality therefore gives variance $O_d(1/p)$. The empirical density has variance at most $1/(4p)$. Replacing δ by the empirical density δ_A changes the centered cube polynomial by at most $2^d|\delta_A - \delta|$. Consequently, for sufficiently large p there is a realization simultaneously approximating the prescribed density, cube count, and centered cube power to any specified positive error.

Apply this to the cosine weights of Proposition III.3.7, with $t_j \rightarrow 0$, and choose primes p_j sufficiently large that all errors are $o(t_j^8)$. Then

$$\alpha_j = t_j^8/32 + o(t_j^8), \quad Q_3(1_{A_j}) - \delta_j^8 = \frac{3}{2}\delta_j^4 t_j^4 + o(t_j^4),$$

which proves (III.3.25).

For the two-frequency example, choose approximation errors $o(u^{28/3})$. The error in replacing its prescribed u^4 by the rounded value is then $o(u^{16/3})$, since u^4 is the square root of the centered eighth power. The density and cube-count errors are smaller as well. Thus (III.3.22) transfers with its nonzero $u^{16/3}$ term. The 3-torsion quintic example can instead be rounded on $\mathbb{Z}/3\mathbb{Z} \times \mathbb{Z}/p\mathbb{Z}$; it does not give quintic sharpness on prime cyclic groups. $\square$

III.4 Density-weighted progression counting

Source 04, delivered Section 7.

The expansion in Gowers's Corollary 3.3 uses every subset of the set of positions, producing a factor bounded by 2^k . An ordered telescoping expansion gives a smaller bound and keeps the density factors. This section first states the precise group hypothesis needed for the Cauchy–Schwarz argument.

Definition III.4.1. Integers $c_1, \dots, c_m$ are *admissible on G* if the map $x \mapsto (c_i - c_j)x$ is an automorphism of G for every $i \neq j$. For functions $a_i : G \rightarrow \mathbb{C}$, write

$$\Lambda_{\mathbf{c}}(a_1, \dots, a_m) = \mathbb{E}_{x,y} \prod_{i=1}^m a_i(x + c_i y).$$

For $G = \mathbb{Z}/N\mathbb{Z}$, admissibility is equivalent to $\gcd(c_i - c_j, N) = 1$ for all pairs. Consecutive slopes $0, 1, \dots, k-1$ are therefore admissible whenever $\gcd(N, (k-1)!) = 1$. Primality is sufficient but unnecessary.

Lemma III.4.2 (Generalized von Neumann bound). *Suppose $m \geq 2$, the slopes are admissible, and $|a_i| \leq 1$. For every j ,*

$$|\Lambda_{\mathbf{c}}(a_1, \dots, a_m)| \leq \|a_j\|_{U^{m-1}}.$$

Proof. Relabel so that the target is a_m . For $m = 2$, the map $(x, y) \mapsto (x + c_1 y, x + c_2 y)$ is a bijection: its inverse uses the automorphism $c_2 - c_1$. The average factors as $(\mathbb{E}a_1)(\mathbb{E}a_2)$, giving the assertion.

Assume $m > 2$ and the result for $m - 1$. Replace x by $x - c_1 y$, and apply Cauchy–Schwarz in x , using $|a_1| \leq 1$. Writing $d_i = c_i - c_1$, this gives

$$|\Lambda_{\mathbf{c}}|^2 \leq \mathbb{E}_x \left| \mathbb{E}_y \prod_{i=2}^m a_i(x + d_i y) \right|^2 = \mathbb{E}_h \mathbb{E}_{x,y} \prod_{i=2}^m a_i(x + d_i y) \overline{a_i(x + d_i y + d_i h)}.$$

For each h , apply the inductive assertion to the $m - 1$ derivative functions. Differences of the new slopes are still automorphisms. The triangle inequality bounds the last expression by

$$\mathbb{E}_h \|\Delta_{d_m h} a_m\|_{U^{m-2}}, \quad \Delta_t a(x) = a(x) \overline{a(x + t)}.$$

The map $h \mapsto d_m h$ is a bijection. Hölder’s inequality and the derivative identity for Gowers norms give

$$\mathbb{E}_h \|\Delta_{d_m h} a_m\|_{U^{m-2}} \leq \left(\mathbb{E}_t \|\Delta_t a_m\|_{U^{m-2}}^{2^{m-2}} \right)^{1/2^{m-2}} = \|a_m\|_{U^{m-1}}^2.$$

Taking square roots completes the induction. $\square$

Theorem III.4.3 (Telescoping with the density factors retained). *Let $k \geq 3$, let the slopes be admissible on G , and let $a_i : G \rightarrow [0, 1]$ have means δ_i . Put $b_i = a_i - \delta_i$ and $u_i = \|b_i\|_{U^{i-1}}$ for $i \geq 3$. Then*

$$\left| \Lambda_{\mathbf{c}}(a_1, \dots, a_k) - \prod_{i=1}^k \delta_i \right| \leq \sum_{j=3}^k \left(\prod_{i=j+1}^k \delta_i \right) u_j. \quad (*)$$

Every permutation of the positions gives a valid version, with the appropriate Gowers norm of the function in the j -th position.

Proof. Successively write $a_j = \delta_j + b_j$, starting with $j = k$ and ending with $j = 3$. Multilinearity gives the exact identity

$$\Lambda_{\mathbf{c}}(a_1, \dots, a_k) - \prod_i \delta_i = \sum_{j=3}^k \left(\prod_{i=j+1}^k \delta_i \right) \Lambda_{c_1, \dots, c_j}(a_1, \dots, a_{j-1}, b_j).$$

The remaining two-function average is exactly $\delta_1 \delta_2$ by pairwise admissibility. Each term is bounded by u_j , by Lemma III.4.2; all other functions in that term are 1-bounded. Taking absolute values proves the result. Relabeling preserves admissibility, proving the last assertion. $\square$

Under the source’s hypotheses $u_j \leq \varepsilon$, the new error is

$$\varepsilon \sum_{j=3}^k \prod_{i=j+1}^k \delta_i \leq (k - 2)\varepsilon.$$

This is a direct replacement for $2^k \varepsilon$, in the same normalization. When all densities equal δ , the factor is

$$S_{k-2}(\delta) := \sum_{r=0}^{k-3} \delta^r = \begin{cases} (1 - \delta^{k-2})/(1 - \delta), & \delta \neq 1, \\ k - 2, & \delta = 1. \end{cases}$$

In particular it is at most $1/(1 - \delta)$ for $0 \leq \delta < 1$, independently of k , and the inequality is strict when $0 < \delta < 1$.

[write] **6 October 2026.** Source 03's Section III.1 proves the same generalized von Neumann bound (Lemma III.1.1; "slopes with invertible differences" there are "admissible slopes" here) and the same telescoping bound (Theorem III.1.2), by the same induction and expansion. Its order-optimization recursion $B(S)$ is $D(S)$ below. The two texts were written independently (the intake measured under 0.8% shared text); both are kept and credited. The first display of Corollary III.4.4 is source 03's (III.1.2). Source 03 concludes at least $\delta^k N^2/4$ progressions under $N \geq 4\delta^{1-k}$ (Corollary III.1.3), while source 04 concludes positivity under $N > 2\delta^{1-k}$. Source 04's exchange rule and source 03's real-character torsion example are each one source's alone. Source 04's $\mathbb{Z}/2$ example below has a complex a_3 and shows that the generalized von Neumann inequality itself fails.

III.4.1 Choosing the order of elimination

The preceding result has a small useful optimization principle. Suppose one has upper bounds $u_i(s) \geq \|b_i\|_{U^s}$, for all $2 \leq s \leq k-1$. For a subset S of positions define

$$D(S) = 0 \quad (|S| \leq 2), \quad D(S) = \min_{i \in S} (u_i(|S| - 1) + \delta_i D(S \setminus \{i\})) \quad (|S| \geq 3).$$

Induction on $|S|$ shows that $D(S)$ bounds the counting error for the subsystem indexed by S . Moreover it is precisely the smallest bound obtained by this family of telescoping estimates. Computing the recurrence requires at most $k2^k$ elementary comparisons, once the norm bounds have been supplied.

If a single upper bound v_i is used for each function at every remaining order, there is also a simple exchange rule. For two successive eliminations i, j , their contribution before the common later tail is either $v_i + \delta_i v_j$ or $v_j + \delta_j v_i$. Thus i should precede j exactly when

$$v_i(1 - \delta_j) \leq v_j(1 - \delta_i).$$

For a fixed pair of positions to retain, sort the other positions by $v_i/(1 - \delta_i)$, with the degenerate cases interpreted through the displayed inequality. One can compare the $\binom{k}{2}$ possible retained pairs. This optimizes the stated bound, and is not a claim of an optimal counting inequality among all methods.

III.4.2 An explicit supersaturation consequence

Corollary III.4.4 (A sharper uniformity threshold). *Let $A \subset G$ have density $0 < \delta \leq 1$, and suppose that the consecutive slopes $0, \dots, k-1$ are admissible, with $k \geq 3$. Put $u = \|1_A - \delta\|_{U^{k-1}}$. The number $T_k^*(A)$ of ordered k -term progressions with nonzero difference satisfies*

$$T_k^*(A) \geq N^2(\delta^k - u S_{k-2}(\delta)) - \delta N. \quad (\text{III.4.1})$$

In particular, if

$$u \leq \frac{\delta^k}{2S_{k-2}(\delta)} \quad \text{and} \quad N > 2\delta^{1-k},$$

then $T_k^(A) > 0$.*

Proof. Norm monotonicity bounds every U^{j-1} -norm in Theorem III.4.3 by u . There are N^2 ordered choices of start and difference. Exactly $|A| = \delta N$ of the counted progressions have difference zero, so subtracting them gives the asserted bound. The two displayed assumptions make its right-hand side positive. $\square$

Admissibility also ensures that a progression with nonzero difference has k distinct entries: a repetition would give $(i - j)y = 0$ with $0 < |i - j| < k$, contrary to the automorphism hypothesis.

[write] **6 October 2026; sources 06–10.** *Two further counting refinements (sources 06 and 10).* Source 10’s Corollary III.29.2 has error $\delta S_k(\delta)u$, a factor δ below the error $uS_{k-2}(\delta)$ of Corollary III.4.4 (the polynomials are equal: $S_k(\delta)$ there is $1 + \delta + \dots + \delta^{k-3}$), because its mixed L^p norms of indicators equal δ ; its Corollary III.29.3 needs $u \leq \delta^{k-1}/(2S_k(\delta))$ in place of $u \leq \delta^k/(2S_{k-2}(\delta))$. Source 06’s Corollary III.5.5 is quadratic in u : error $\sqrt{\delta} B_k(\delta)u^2$ with $B_k(\delta) = \sum_{t \leq k-3} (t+1)\delta^t$, and its Corollary III.5.7 needs $u^2 \leq \delta^{k-1/2}/(2B_k(\delta))$ (corrected after the independent check of 6 October 2026; the first wording, “its criterion”, read as a reference to Corollary III.5.5). In the paper’s α (with $u^{2^{k-1}} \leq \alpha$) the sufficient conditions become $\alpha \leq (\delta^k/2S_{k-2})^{2^{k-1}}$ here, $\alpha \leq (\delta^{k-1}/2S_k)^{2^{k-1}}$ (source 10) and $\alpha \leq (\delta^{k-1/2}/2B_k)^{2^{k-2}}$ (source 06), against the formal `corollary_3_6`’s $(\delta/2)^{k2^k}$. The powers of δ are $k2^{k-1}$, $(k-1)2^{k-1}$, $(k-\frac{1}{2})2^{k-2}$ and $k2^k$; source 06’s is the smallest for every $k \geq 3$, since $(k-\frac{1}{2})2^{k-2} < (k-1)2^{k-1}$ is $k > \frac{3}{2}$ (the constants $2S_{k-2} \leq 2(k-2)$ and $2B_k \leq (k-1)(k-2)$ differ). All need invertibility of $1, \dots, k-1$; source 10’s Theorem III.28.2 and its criterion (III.29.6) work without it, with explicit kernel factors.

In the paper’s uniformity parameter, the sufficient condition becomes

$$\alpha \leq \left(\frac{\delta^k}{2S_{k-2}(\delta)} \right)^{2^{k-1}}.$$

For fixed k and small δ , its density exponent is $k2^{k-1}$, compared with $k2^k$ in the sufficient condition $(\delta/2)^{k2^k}$ displayed in the inspected version of Corollary 3.6. This comparison concerns the modular progression step.

Lemma III.4.5 (Lifting from a short interval). *Let $A \subset \{0, \dots, n-1\}$ be viewed in $\mathbb{Z}/N\mathbb{Z}$, with $2(n-1) < N$. If A contains a nonconstant modular progression of length $k \geq 3$, then it contains an ordinary integer progression of length k and positive difference.*

Proof. Let $x_0, \dots, x_{k-1} \in \{0, \dots, n-1\}$ be the representatives. Each integer $x_i + x_{i+2} - 2x_{i+1}$ is divisible by N and has absolute value at most $2(n-1) < N$; it is therefore zero. Successive integer differences are equal. The common difference is nonzero because the modular progression is nonconstant. Reverse the sequence if that integer difference is negative. $\square$

The lemma specifies the extra hypothesis needed to draw an integer conclusion. A modular progression in an unrestricted subset of $\mathbb{Z}/N\mathbb{Z}$ can wrap around, so that conclusion should not be silently inserted into Corollary III.4.4.

[write] **6 October 2026.** *Already formalized.* This lemma is proved in Lean as `hasNatAP_of_short_modular_sequence` (`Proofs18IntervalTransfer.lean`, line 22), with the set form `hasNatAP_of_hasModAP_image` (line 92); both were added in commit 327773149, source 04’s own pin. The Lean statement takes representatives $f(i) \leq n$ with $2n < N$ and a nonzero modular difference, for every length k : after $n-1 \mapsto n$ it is this lemma, and it also covers $k < 3$. Source 04 does not mention the formal version. It is the only statement of sources 01–05 that is formalized; source 03’s rectification remark (Remark III.1.4) is the same argument.

Remark III.4.6 (Why invertibility cannot simply be omitted). For $G = \mathbb{Z}/2\mathbb{Z}$, take slopes 0, 1, 2, $a_3(0) = 1$, $a_3(1) = i$, $a_1 = \overline{a_3}$, and $a_2 = 1$. The forms at positions 0 and 2 coincide, so $\Lambda = 1$. Both normalized Fourier coefficients of a_3 have modulus $2^{-1/2}$, whence $\|a_3\|_{U^2} = 2^{-1/4} < 1$. Thus the stated generalized von Neumann inequality itself fails without an algebraic hypothesis. The obstruction is an algebraic dependence, and the admissibility condition states exactly what the present proof uses.

III.5 Quadratic discrepancy bounds for arithmetic progressions

Source 06, delivered Section 3.

The proof of Gowers's Theorem 3.2 is a repeated Cauchy–Schwarz argument. Keeping two functions in that argument, rather than discarding all but one of their norms, produces a quadratic error estimate. A second telescoping identity then exploits the exact disappearance of one- and two-function balanced correlations. This refines the displayed counting estimates in Section 3 of the paper and its repository transcription [1, 41].

III.5.1 Normalization and the precise hypothesis on the modulus

Let G be a finite abelian group of order N , with normalized averages. For an integer a , let $[a]: G \rightarrow G$ denote multiplication by a . For slopes $a_1, \dots, a_k$, define

$$\Lambda_a(F_1, \dots, F_k) = \mathbb{E}_{x,d \in G} \prod_{i=1}^k F_i(x + a_i d).$$

Throughout this section, assume that

$$[a_i - a_j] \text{ is an automorphism of } G \text{ whenever } i \neq j. \quad (\text{III.5.1})$$

This assumption is inherited by every subset of the slopes. It makes the map

$$(x, d) \mapsto (x + a_i d, x + a_j d)$$

a bijection of G^2 . In particular,

$$\mathbb{E}_{x,d} F_i(x + a_i d) F_j(x + a_j d) = (\mathbb{E} F_i)(\mathbb{E} F_j). \quad (\text{III.5.2})$$

For consecutive slopes $0, 1, \dots, k-1$, condition (III.5.1) means that multiplication by each of $1, \dots, k-1$ is invertible. For $G = \mathbb{Z}/N\mathbb{Z}$ it is equivalent to

$$\gcd(N, (k-1)!) = 1. \quad (\text{III.5.3})$$

Thus the counting results below apply to many composite moduli. For example, $N = 25$ is admissible for every $k \leq 5$. If $N \geq 2$ and the slopes are consecutive, every progression with $d \neq 0$ has k distinct terms: an equality between two terms would put d in the kernel of one of the automorphisms in (III.5.1).

For $s \geq 1$, use the normalized uniformity quantity

$$\|f\|_{U^s(G)}^{2^s} = \mathbb{E}_{x, h_1, \dots, h_s \in G} \prod_{\omega \in \{0,1\}^s} \mathcal{C}^{|\omega|} f(x + \omega \cdot h),$$

where $\mathcal{C}$ is complex conjugation. Thus $\|f\|_{U^1} = |\mathbb{E} f|$. The standard derivative identity is

$$\|f\|_{U^s}^{2^s} = \mathbb{E}_h \|\Delta_h f\|_{U^{s-1}}^{2^{s-1}}, \quad \Delta_h f(x) = f(x) \overline{f(x+h)} \quad (s \geq 2). \quad (\text{III.5.4})$$

We use the usual inequalities $\|f\|_{U^s} \leq \|f\|_{U^{s+1}}$ and $\|f\|_{U^s} \leq \|f\|_\infty$; they follow from the same Cauchy–Schwarz argument used in Section 3 of [1]. In the paper's convention, degree- d α -uniformity is precisely

$$\|f\|_{U^{d+1}}^{2^{d+1}} \leq \alpha. \quad (\text{III.5.5})$$

III.5.2 Two distinguished functions

Theorem III.5.1 (Two-function uniformity estimate). *Assume (III.5.1) and $k \geq 2$. For distinct indices p, q and arbitrary complex functions $F_1, \dots, F_k$ on G ,*

$$|\Lambda_a(F_1, \dots, F_k)| \leq \|F_p\|_{U^{k-1}} \|F_q\|_{U^{k-1}} \prod_{i \notin \{p, q\}} \|F_i\|_\infty. \quad (\text{III.5.6})$$

If $k \geq 3$ and $r \notin \{p, q\}$, then the stronger estimate

$$|\Lambda_a(F_1, \dots, F_k)| \leq \|F_r\|_2 \|F_p\|_{U^{k-1}} \|F_q\|_{U^{k-1}} \prod_{i \notin \{p, q, r\}} \|F_i\|_\infty \quad (\text{III.5.7})$$

also holds. The coefficient 1 in (III.5.6) is sharp.

Proof. By homogeneity, it is enough to suppose that the functions other than F_p, F_q have supremum norm at most one. If any normalizing norm is zero, the corresponding function vanishes and there is nothing to prove. For $k = 2$, equation (III.5.2) gives the identity

$$|\Lambda_{(a_p, a_q)}(F_p, F_q)| = |\mathbb{E} F_p| |\mathbb{E} F_q| = \|F_p\|_{U^1} \|F_q\|_{U^1}.$$

Suppose (III.5.6) has been proved for $k - 1$ functions. Choose $r \notin \{p, q\}$ and change x to $x - a_r d$ so that the r th slope becomes zero. Put $b_i = a_i - a_r$. Cauchy–Schwarz in x gives

$$\begin{aligned} |\Lambda_a(F_1, \dots, F_k)|^2 &\leq \|F_r\|_2^2 \mathbb{E}_x \left| \mathbb{E}_d \prod_{i \neq r} F_i(x + b_i d) \right|^2 \\ &= \|F_r\|_2^2 \mathbb{E}_h \Lambda_{(b_i)_{i \neq r}} ((\Delta_{b_i h} F_i)_{i \neq r}). \end{aligned}$$

The last average is real and nonnegative, as is apparent from the preceding line. Applying the triangle inequality and the induction hypothesis bounds it by

$$\mathbb{E}_h B(h) C(h), \quad B(h) = \|\Delta_{b_p h} F_p\|_{U^{k-2}}, \quad C(h) = \|\Delta_{b_q h} F_q\|_{U^{k-2}}.$$

All remaining differentiated functions have supremum norm at most one. Set $Q = 2^{k-2} \geq 2$. Hölder’s inequality on the probability space G gives

$$\mathbb{E}_h B(h) C(h) \leq (\mathbb{E}_h B(h)^Q)^{1/Q} (\mathbb{E}_h C(h)^Q)^{1/Q}.$$

For $Q > 2$, this is Hölder with a third factor 1 and exponent $Q/(Q - 2)$; for $Q = 2$, it is Cauchy–Schwarz. Both $[b_p]$ and $[b_q]$ are automorphisms. Therefore (III.5.4) yields

$$\mathbb{E}_h B(h)^Q = \|F_p\|_{U^{k-1}}^{2^{k-1}}, \quad \mathbb{E}_h C(h)^Q = \|F_q\|_{U^{k-1}}^{2^{k-1}}.$$

Taking square roots proves (III.5.7). Using $\|F_r\|_2 \leq \|F_r\|_\infty$ proves (III.5.6) and completes the induction. Restoring the supremum norms gives the homogeneous statements. Constant functions show that the coefficient in (III.5.6) cannot be decreased. $\square$

[write] **6 October 2026; sources 15–20.** Source 20’s Theorem III.43.3 proves (III.5.7) for arbitrary surjective homomorphisms with directional seminorms, independently (source 20 could not see this manuscript).

Bounding one of the two distinguished uniformity quantities by the supremum norm recovers the ordinary generalized von Neumann estimate. The gain in Theorem III.5.1 is useful when both distinguished functions are balanced and have small uniformity norm.

[write] **6 October 2026; sources 06–10.** The one-function case is the generalized von Neumann bound of sources 03 and 04 (Lemmas III.1.1 and III.4.2), with the same slope hypothesis (III.5.1). Source 10's Theorem III.28.2 replaces the L^∞ norms by mixed L^p norms and drops the invertibility of all differences except one, at the price of an explicit kernel factor. Keeping two distinguished functions is specific to this source.

III.5.3 First and second telescoping

Let $0 \leq F_i \leq 1$, let $\delta_i = \mathbb{E}F_i$, and put $f_i = F_i - \delta_i$.

Theorem III.5.2 (Indexed first-order estimate). *Suppose $k \geq 2$, assume (III.5.1), and suppose $\|f_j\|_{U^{j-1}} \leq \varepsilon_j$ for $3 \leq j \leq k$. Then*

$$\left| \Lambda_a(F_1, \dots, F_k) - \prod_{i=1}^k \delta_i \right| \leq \sum_{j=3}^k \varepsilon_j \prod_{h=j+1}^k \delta_h. \quad (\text{III.5.8})$$

For $k = 2$, the empty sum is zero and the count is exact.

Proof. The pointwise telescoping identity is

$$\prod_{i=1}^k F_i - \prod_{i=1}^k \delta_i = \sum_{j=1}^k f_j \prod_{i < j} F_i \prod_{h > j} \delta_h.$$

Evaluate the i th function at $x + a_i d$ and average. The $j = 1$ summand vanishes because $\mathbb{E}f_1 = 0$. The $j = 2$ summand vanishes by (III.5.2). For $j \geq 3$, apply the one-function consequence of Theorem III.5.1 to the j slopes $a_1, \dots, a_j$. This bounds that summand by $\varepsilon_j \prod_{h > j} \delta_h$. $\square$

Under the indexed hypotheses of Gowers's Corollary 3.3, the right side of (III.5.8) is at most $(k-2)\alpha$, in the notation of that corollary, replacing its displayed $2^k\alpha$ bound. The symbol α there is a norm-scale parameter; conversion to the paper's degree-uniformity energy must use (III.5.5).

Theorem III.5.3 (Second-order estimate). *Suppose $k \geq 3$, assume (III.5.1), and suppose $\|f_i\|_{U^{k-1}} \leq \varepsilon_i$ for all i . For $2 \leq i < j \leq k$, set*

$$w_{ij} = \prod_{i < h < j} \delta_h \prod_{h > j} \delta_h, \quad \mu_i = \min_{r < i} \|F_r\|_2.$$

Then

$$\left| \Lambda_a(F_1, \dots, F_k) - \prod_{i=1}^k \delta_i \right| \leq \sum_{2 \leq i < j \leq k} w_{ij} \mu_i \varepsilon_i \varepsilon_j. \quad (\text{III.5.9})$$

One may replace every μ_i by 1. For indicator functions, $\mu_i = \min_{r < i} \sqrt{\delta_r}$.

Proof. In the proof of Theorem III.5.2, telescope each prefix a second time:

$$\prod_{h < j} F_h = \prod_{h < j} \delta_h + \sum_{i < j} f_i \prod_{r < i} F_r \prod_{i < h < j} \delta_h.$$

The constant-prefix contribution vanishes after multiplication by $f_j(x + a_j d)$ and averaging. The $i = 1$ contribution vanishes by (III.5.2). The exact surviving expression for the discrepancy is consequently

$$\sum_{2 \leq i < j \leq k} w_{ij} \Lambda_{(a_1, \dots, a_i, a_j)}(F_1, \dots, F_{i-1}, f_i, f_j).$$

There are $i + 1$ functions in each average. Apply (III.5.7) with f_i, f_j distinguished and any F_r , $r < i$, supplying the L^2 factor. Norm monotonicity gives $\|f_i\|_{U^i} \leq \|f_i\|_{U^{k-1}}$ and the corresponding inequality for f_j . Minimizing over $r < i$ and summing proves the claim. $\square$

Remark III.5.4 (Ordering the inputs). If all balanced functions satisfy the common bound $\|f_i\|_{U^{k-1}} \leq \varepsilon$, the slope–function pairs may be reordered before either telescoping argument. A nonincreasing order of the densities minimizes the coefficient in (III.5.8), and also minimizes the coefficient in (III.5.9) when the μ_i are replaced by 1. For the first assertion, the coefficient is a sum of products of final segments. In an adjacent swap, the only changed final-segment term is a nonnegative common product times the later density; placing the smaller density later cannot increase it. For the second assertion, the coefficient is

$$\sum_{i=2}^{k-1} e_{k-i-1}(\delta_{i+1}, \dots, \delta_k),$$

where e_j is the elementary symmetric polynomial. Again, an adjacent swap changes only one final-segment summand, whose dependence on the changed endpoint is affine with a nonnegative coefficient. This is an optimization of these bounds among orders; it does not assert that their coefficients are universally sharp discrepancy constants.

III.5.4 One set and explicit coefficients

For $k \geq 3$, define the polynomials

$$S_k(\delta) = \sum_{t=0}^{k-3} \delta^t, \quad B_k(\delta) = \sum_{t=0}^{k-3} (t+1) \delta^t. \quad (\text{III.5.10})$$

Corollary III.5.5 (Quadratic progression discrepancy). *Let $A \subseteq G$ have density δ , let $\varepsilon = \|1_A - \delta\|_{U^{k-1}}$, and suppose $k \geq 3$ and (III.5.1). Then*

$$\left| \Lambda_a(1_A, \dots, 1_A) - \delta^k \right| \leq \min \left\{ S_k(\delta) \varepsilon, \sqrt{\delta} B_k(\delta) \varepsilon^2 \right\}. \quad (\text{III.5.11})$$

Proof. Apply Theorems III.5.2 and III.5.3, using norm monotonicity for the first. In the second theorem, $\mu_i = \sqrt{\delta}$ and $w_{ij} = \delta^{k-i-1}$. For fixed i there are $k - i$ choices of j , so its coefficient is

$$\sum_{i=2}^{k-1} (k-i) \delta^{k-i-1} = \sum_{t=0}^{k-3} (t+1) \delta^t = B_k(\delta).$$

$\square$

[write] **6 October 2026; sources 15–20.** Source 20’s Corollary III.45.4 has a coefficient (its own $B_k(\delta)$, a different polynomial) at most $\sqrt{\delta} B_k(\delta)$ with the B_k here, and strictly smaller for every $k \geq 3$ and $0 < \delta < 1$ (proved in the note there); the quadratic power itself is proved here first, and source 20’s Theorem III.48.3 shows it cannot be raised for three-term counts of indicator sets at density near $1/2$.

For $0 \leq \delta < 1$, the closed forms are

$$S_k(\delta) = \frac{1 - \delta^{k-2}}{1 - \delta}, \quad B_k(\delta) = \frac{1 - (k-1)\delta^{k-2} + (k-2)\delta^{k-1}}{(1 - \delta)^2}.$$

At $\delta = 1$, the polynomial values are $k - 2$ and $(k - 1)(k - 2)/2$. In particular,

$$B_k(\delta) \leq \frac{(k-1)(k-2)}{2}, \quad B_k(\delta) \leq \frac{1}{(1-\delta)^2} \quad (\delta < 1).$$

For fixed $\delta < 1$, this latter bound remains bounded as k grows. The quadratic dependence on ε strengthens the usual linear discrepancy estimate independently of the improvement in the coefficient.

Proposition III.5.6 (A sharp three-term variance factor). *Under the three-slope case of (III.5.1),*

$$\left| \Lambda_a(1_A, 1_A, 1_A) - \delta^3 \right| \leq \sqrt{\delta(1-\delta)} \|1_A - \delta\|_{U^2}^2. \quad (\text{III.5.12})$$

The coefficient 1 is sharp, even for consecutive slopes in cyclic groups of odd order.

Proof. Set $f = 1_A - \delta$. Every expansion term with one or two factors f vanishes by mean zero and (III.5.2). Thus the discrepancy equals $\Lambda_a(f, f, f)$. Apply (III.5.7) and use $\|f\|_2^2 = \delta(1 - \delta)$.

For sharpness, take A to be a singleton in a cyclic group of odd order $N \geq 3$. Then $\delta = 1/N$, and its normalized three-term count is $1/N^2$. With normalized Fourier coefficients, the balanced singleton has coefficient 0 at the trivial character and modulus $1/N$ at each of the other $N - 1$ characters. Therefore

$$\|f\|_{U^2}^4 = \frac{N-1}{N^4}, \quad \sqrt{\delta(1-\delta)} \|f\|_{U^2}^2 = \frac{N-1}{N^3} = \frac{1}{N^2} - \frac{1}{N^3}.$$

Equality holds in (III.5.12). □

[write] **6 October 2026; sources 15–20.** This coefficient is source 20's $B_3(\delta)$. Its sharpness here is at singletons ($\delta = 1/N$); at fixed density source 20 shows that the exponent 2 cannot be raised (Theorem III.48.3), while the best coefficient stays open (item III.8).

III.5.5 An explicit modular existence criterion

Corollary III.5.7 (Exact subtraction of constant progressions). *Let $N \geq 2$, $k \geq 3$, and let $A \subseteq G$ have density $0 < \delta \leq 1$. Assume that multiplication by $1, \dots, k-1$ is invertible on G . Put*

$$E_k = \min \left\{ S_k(\delta)\varepsilon, \sqrt{\delta} B_k(\delta)\varepsilon^2 \right\}, \quad \varepsilon = \|1_A - \delta\|_{U^{k-1}}.$$

The number of nonconstant ordered k -term progressions in A is at least

$$N^2(\delta^k - E_k) - \delta N. \quad (\text{III.5.13})$$

In particular, such a progression exists whenever

$$\delta^k - E_k > \frac{\delta}{N}. \quad (\text{III.5.14})$$

A convenient sufficient condition is

$$N > 2\delta^{1-k}, \quad \varepsilon^2 \leq \frac{\delta^{k-1/2}}{2B_k(\delta)}. \quad (\text{III.5.15})$$

Consequently, if A is α -uniform of degree $k-2$ in the paper's convention, it is sufficient that

$$N > 2\delta^{1-k}, \quad \alpha \leq \left(\frac{\delta^{k-1/2}}{2B_k(\delta)} \right)^{2^{k-2}}. \quad (\text{III.5.16})$$

Proof. Corollary III.5.5 gives at least $N^2(\delta^k - E_k)$ ordered progressions. Exactly $|A| = \delta N$ of them have $d = 0$. Subtracting these proves (III.5.13); every remaining progression has distinct terms by (III.5.1). Under (III.5.15), $E_k \leq \delta^k/2$ and $\delta/N < \delta^k/2$, proving (III.5.14). Finally, (III.5.5) implies $\varepsilon^2 \leq \alpha^{1/2^{k-2}}$, giving (III.5.16). $\square$

For $k = 2$, no uniformity assumption is necessary: the number of nonconstant progressions is exactly $|A|(|A| - 1)$, so existence is equivalent to $|A| \geq 2$. For $\delta = 1$ and $N \geq 2$, a nonzero difference gives a progression immediately. For $N = 1$, no nonconstant progression exists. Empty and full sets have zero balanced function, so their progression counts are exact.

The conclusions of Corollary III.5.7 concern modular progressions. Transferring them to integer intervals requires a separate embedding without wraparound or a rectification argument. The improved uniformity threshold is a local improvement in the counting part of the proof; a stronger final quantitative Szemerédi theorem would also require corresponding control of the subsequent inverse and density-increment steps.

[write] **6 October 2026; sources 06–10.** Theorem III.5.2 is the telescope of sources 03 and 04 (Theorems III.1.2 and III.4.3); Remark III.5.4 is their optimal-order observation. The criteria of Corollary III.5.7 are compared with those of sources 03, 04 and 10 in the note after Corollary III.4.4: in the paper's α , (III.5.16) has the smallest power of δ . Not formalized; `corollary_3_3_holds` and `corollary_3_6_holds` keep the printed constants.

III.6 Normalization and the basic analytic tools

Source 08, delivered Section 2.

Let G be a finite abelian group, written additively, and let all expectations be normalized counting averages. Put

$$V_d = \{0, 1\}^d, \quad p = 2^d, \quad L_v(x, h) = x + \sum_{i=1}^d v_i h_i.$$

The variables $(x, h_1, \dots, h_d)$ are independent and uniform in G . Define the real multilinear cube average by

$$\Lambda_d((F_v)_{v \in V_d}) = \mathbb{E}_{x, h} \prod_{v \in V_d} F_v(L_v(x, h)), \quad \mathcal{C}_d(F) = \Lambda_d((F)_{v \in V_d}). \quad (\text{III.6.1})$$

For a set $A \subseteq G$, the unnormalized number of parameterized cubes in A is $|G|^{d+1} \mathcal{C}_d(\mathbf{1}_A)$. Degenerate cubes are included. Thus no assumption that the sampled vertices are distinct is hidden in the notation.

For a complex function g , the normalized Gowers norm is

$$\|g\|_{U^d(G)}^p = \mathbb{E}_{x,h} \prod_{v \in V_d} \mathcal{K}^{|v|} g(L_v(x, h)), \quad (\text{III.6.2})$$

where $\mathcal{K}$ is complex conjugation. For real functions,

$$\mathcal{C}_d(g) = \|g\|_{U^d(G)}^p. \quad (\text{III.6.3})$$

For $d = 1$ this is a seminorm, $\|g\|_{U^1} = |\mathbb{E}g|$; for $d \geq 2$ it is a norm.

We use the normalized Gowers–Cauchy–Schwarz inequality

$$\left| \mathbb{E}_{x,h} \prod_{v \in V_d} \mathcal{K}^{|v|} g_v(L_v) \right| \leq \prod_{v \in V_d} \|g_v\|_{U^d(G)}. \quad (\text{III.6.4})$$

This is the normalized form of the cube inequality in [1, Lemma 3.8]. Its repeated Cauchy–Schwarz proof uses only translation invariance of the uniform measure, so it applies to every finite abelian group. Real inputs allow us to omit the conjugations. Complex unconjugated inputs can be treated by replacing g_v with $\mathcal{K}^{|v|} g_v$.

Here are the consequences needed later. Padding one face of a $(d+1)$ -cube with the constant function 1, and the other face with g , gives from (III.6.4)

$$\|g\|_{U^d}^{2^d} \leq \|g\|_{U^{d+1}}^{2^d}.$$

Together with the U^1 identity this proves

$$|\mathbb{E}g| \leq \|g\|_{U^2} \leq \cdots \leq \|g\|_{U^d} \leq \|g\|_{\infty}. \quad (\text{III.6.5})$$

Expanding the cube of $g + h$ and applying (III.6.4) term by term proves Minkowski’s inequality. These facts require no condition on the torsion of G .

Fix now a real function F with nonnegative mean, and write

$$\delta = \mathbb{E}F \geq 0, \quad f = F - \delta, \quad t = \|f\|_{U^d(G)}, \quad \mathcal{D}_d(F) = \mathcal{C}_d(F) - \delta^p. \quad (\text{III.6.6})$$

Most applications use $0 \leq F \leq 1$, but the polynomial identities hold for arbitrary real f of mean zero and arbitrary $\delta \geq 0$. By (III.6.5), $\mathcal{D}_d(F) \geq 0$ even without pointwise positivity of F .

For $S \subseteq V_d$, set

$$\Lambda_S(f) = \mathbb{E}_{x,h} \prod_{v \in S} f(L_v(x, h)), \quad \Lambda_{\emptyset}(f) = 1. \quad (\text{III.6.7})$$

Padding the other vertices with 1 in (III.6.4) gives the useful homogeneous estimate

$$|\Lambda_S(f)| \leq t^{|S|}. \quad (\text{III.6.8})$$

Unlike the estimate of every error term by t , this remembers the number of balanced factors.

III.7 Three exact cancellations and the fourth-order bound

Source 08, delivered Section 3.

Lemma III.7.1 (Three-vertex independence). *Let $1 \leq r \leq 3$, and let $v_1, \dots, v_r$ be distinct vertices of V_d . Then $(L_{v_1}, \dots, L_{v_r})$ is uniformly distributed on G^r . Consequently, for functions $g_1, \dots, g_r : G \rightarrow \mathbb{C}$,*

$$\mathbb{E}_{x,h} \prod_{j=1}^r g_j(L_{v_j}) = \prod_{j=1}^r \mathbb{E} g_j.$$

Proof. Reflecting coordinate i of a cube is implemented by $x' = x + h_i$, $h'_i = -h_i$, leaving the other increments unchanged. This is a bijection of G^{d+1} preserving its uniform distribution. Reflect enough coordinates to send v_1 to 0.

For two vertices the other vertex is a nonzero binary vector u . At least one coefficient of $u \cdot h$ is 1, so this linear form is uniform and independent of x .

For three vertices write them as $0, u, v$, with u, v distinct and nonzero. There are two coordinates i, j for which

$$\det \begin{pmatrix} u_i & u_j \\ v_i & v_j \end{pmatrix} = \pm 1.$$

Indeed, distinct nonzero binary vectors are linearly independent over $\mathbb{F}_2$; an invertible 2×2 binary minor has integer determinant ± 1 . After fixing all other increments, the map $(h_i, h_j) \mapsto (u \cdot h, v \cdot h)$ is therefore a bijection of G^2 : its inverse has integer coefficients. Together with the independent variable x , this makes $(x, x + u \cdot h, x + v \cdot h)$ uniform on G^3 . No division by 2 or another potentially noninvertible integer occurs. $\square$

Corollary III.7.2. *If $\mathbb{E}f = 0$ and $1 \leq |S| \leq 3$, then $\Lambda_S(f) = 0$. The same conclusion holds when the centered function at each selected vertex is allowed to be different.*

Theorem III.7.3 (Fourth-order refinement). *Let G be a finite abelian group, $d \geq 2$, $p = 2^d$, and let $F = \delta + f$ be real, where $\delta \geq 0$ and $\mathbb{E}f = 0$. With $t = \|f\|_{U^d(G)}$,*

$$0 \leq \mathcal{D}_d(F) \leq \sum_{r=4}^p \binom{p}{r} \delta^{p-r} t^r, \quad (\text{III.7.1})$$

$$\mathcal{C}_d(F) \leq (\delta + t)^p - p\delta^{p-1}t - \binom{p}{2} \delta^{p-2} t^2 - \binom{p}{3} \delta^{p-3} t^3. \quad (\text{III.7.2})$$

In dimension 2, the first estimate is an identity: $\mathcal{C}_2(F) = \delta^4 + \|f\|_{U^2}^4$.

Proof. The finite product expansion gives the exact identity

$$\mathcal{C}_d(\delta + f) = \sum_{S \subseteq V_d} \delta^{p-|S|} \Lambda_S(f). \quad (\text{III.7.3})$$

The empty subset contributes δ^p . Corollary III.7.2 removes all subsets of sizes 1, 2, 3. For each remaining subset, (III.6.8) bounds its absolute contribution by $\delta^{p-|S|} t^{|S|}$. There are $\binom{p}{r}$ subsets of size r . This proves the upper bound; the lower bound was established in (III.6.5). Expanding $(\delta + t)^p$ gives (III.7.2). When $d = 2$, the only nonempty surviving subset is the entire four-vertex cube, whose average is t^4 . $\square$

[write] **6 October 2026; sources 06–10.** The cancellation through degree three and (III.7.1) are also proved by source 03 (Lemma III.2.1, Corollary III.2.2), source 04 (Lemma III.3.1) and source 09 (Lemma III.17.1, Theorem III.17.2); four independent proofs. This t is 03's η , 04's u and 09's ε ; this $\mathcal{C}_d$ is their Q_d and $\mathcal{C}_d$. Not formalized; lemma_3_10_holds keeps the Minkowski bound.

III.7.1 Conversion to the original uniformity parameter

Define

$$K_d = \sum_{r=4}^p \binom{p}{r} = 2^p - 1 - p - \binom{p}{2} - \binom{p}{3}. \quad (\text{III.7.4})$$

For $F = \mathbf{1}_A$, $0 \leq \delta, t \leq 1$. If $t^p \leq \alpha \leq 1$, then

$$0 \leq \mathcal{C}_d(\mathbf{1}_A) - \delta^p \leq K_d \alpha^{4/p}. \quad (\text{III.7.5})$$

The original bound (F.13.1), after normalization and subtraction of δ^p , has a leading term of order $\alpha^{1/p}$ for fixed positive δ . Equation (III.7.5) replaces that exponent by $4/p$. For example, the exponents are 1, $1/2$, $1/4$, and $1/8$ in dimensions 2, 3, 4, 5, respectively. Theorem III.11.4 will show that this power cannot be increased in a bound uniform over groups and sets.

III.7.2 Density-sensitive versions

When $0 < t \leq \delta$, equation (III.7.1) gives

$$\mathcal{D}_d(F) \leq K_d \delta^{p-4} t^4. \quad (\text{III.7.6})$$

More precisely, for $\delta > 0$ put $s = t/\delta$. Then

$$\frac{\mathcal{D}_d(F)}{\delta^p} \leq B_p(s), \quad B_p(s) = \sum_{r=4}^p \binom{p}{r} s^r. \quad (\text{III.7.7})$$

This makes the relative scale t/δ , rather than t alone, explicit. For $0 < \varepsilon \leq 1$, a sufficient condition for relative error at most ε is

$$t^p \leq \delta^p \left(\frac{\varepsilon}{K_d} \right)^{p/4}. \quad (\text{III.7.8})$$

The condition implies $t \leq \delta$ because $K_d \geq 1$.

A useful closed bound, valid without $t \leq \delta$, is

$$\mathcal{D}_d(F) \leq \binom{p}{4} t^4 (\delta + t)^{p-4}. \quad (\text{III.7.9})$$

To verify it, compare coefficients using $\binom{p}{4} \binom{p-4}{r-4} = \binom{p}{r} \binom{r}{4} \geq \binom{p}{r}$. These convenient constants are not claimed to be optimal. The next section substantially reduces the fourth-order coefficient itself.

III.8 The exact quartic term and the role of two-torsion

Source 08, delivered Section 4.

III.8.1 Fourier notation

Write $\widehat{G}$ additively for the character group, and use $e_\chi(x)$ for the character corresponding to $\chi \in \widehat{G}$. Set

$$\begin{aligned}\widehat{f}(\chi) &= \mathbb{E}_x f(x) \overline{e_\chi(x)}, & E_4(f) &= \sum_{\chi \in \widehat{G}} |\widehat{f}(\chi)|^4, \\ E_{\langle 2 \rangle}(f) &= \sum_{\substack{\chi \in \widehat{G} \setminus \{0\} \\ 2\chi=0}} |\widehat{f}(\chi)|^4.\end{aligned}\tag{III.8.1}$$

For real f , $\widehat{f}(-\chi) = \overline{\widehat{f}(\chi)}$. Character orthogonality on a two-cube gives

$$E_4(f) = \|f\|_{U^2(G)}^4, \quad 0 \leq E_{\langle 2 \rangle}(f) \leq E_4(f) \leq t^4.\tag{III.8.2}$$

If G has odd order, $E_{\langle 2 \rangle}(f) = 0$. If G has exponent 2 and f has mean zero, $E_{\langle 2 \rangle}(f) = E_4(f)$.

For any $S \subseteq V_d$, expanding every factor by Fourier inversion yields

$$\Lambda_S(f) = \sum_{\substack{(\chi_v)_{v \in S} \in \widehat{G}^S \\ \sum_v \chi_v = 0, \sum_v v_i \chi_v = 0 \ (1 \leq i \leq d)}} \prod_{v \in S} \widehat{f}(\chi_v).\tag{III.8.3}$$

This formula is exact over every finite abelian group. In particular, one must solve the equations in $\widehat{G}$, not merely compute a rank over $\mathbb{Q}$.

III.8.2 Classification of four-vertex subsets

A set of four distinct binary vertices is an *ordinary parallelogram* if it can be labeled a, b, c, e with $a + b = c + e$ as an equality in $\mathbb{Z}^d$. It is a *binary affine plane* if its four vertices sum to 0 in $\mathbb{F}_2^d$. Every ordinary parallelogram is a binary affine plane, but the converse fails in dimension at least 3.

Lemma III.8.1 (Quartic classification). *Suppose $f : G \rightarrow \mathbb{R}$ has mean zero and $|S| = 4$. Then*

$$\Lambda_S(f) = \begin{cases} E_4(f), & S \text{ is an ordinary parallelogram,} \\ E_{\langle 2 \rangle}(f), & S \text{ is a binary affine plane but not an ordinary parallelogram,} \\ 0, & S \text{ is not a binary affine plane.} \end{cases}$$

Proof. Use cube reflections to write $S = \{0, u, v, w\}$ with distinct nonzero binary vectors u, v, w . Let the corresponding Fourier frequencies be $\lambda_0, \lambda_1, \lambda_2, \lambda_3$. The constraints in (III.8.3) are

$$\lambda_0 + \lambda_1 + \lambda_2 + \lambda_3 = 0, \quad u_i \lambda_1 + v_i \lambda_2 + w_i \lambda_3 = 0 \quad (1 \leq i \leq d).\tag{III.8.4}$$

First suppose the matrix with columns u, v, w has rank 2 over $\mathbb{Q}$. The vectors u, v have a unit 2×2 minor, and $w = au + bv$ for some rational a, b . Among the coordinate patterns (u_i, v_i) , at least two of 10, 01, 11 occur. If all three occur, the constraints that $a, b, a + b$ lie in $\{0, 1\}$ would force w to be 0, u , or v , which is excluded. If only 10, 01 occur, then $w = u + v$; if the two patterns are 10, 11 or 01, 11, then one of u, v, w is the sum of the other two. In every case the four points form a parallelogram. The unit minor shows that all solutions of (III.8.4) are the primitive parallelogram relation multiplied by a single character χ . The product of the four Fourier coefficients is $|\widehat{f}(\chi)|^4$.

Next suppose the matrix has rank 3 over $\mathbb{F}_2$. It has an odd 3×3 binary minor. Every 3×3 binary determinant has absolute value at most 2, so this minor is ± 1 and is invertible over every abelian group. Thus all four frequencies vanish, and the contribution is zero because $\hat{f}(0) = 0$. For completeness, the determinant bound follows by adjoining an all-one first row and column and replacing the binary block A by $1 - 2A$. The resulting 4×4 sign matrix has determinant $-8 \det A$ and absolute determinant at most 16 by Hadamard's inequality.

The remaining case has rank 3 over $\mathbb{Q}$ and rank 2 over $\mathbb{F}_2$. Since u, v, w are distinct and nonzero, their binary relation is $u + v + w = 0$. The only nonzero coordinate patterns (u_i, v_i, w_i) are 110, 101, 011. All three patterns must occur, since otherwise the rational rank would be at most 2. Equations (III.8.4) now give

$$\lambda_1 + \lambda_2 = \lambda_1 + \lambda_3 = \lambda_2 + \lambda_3 = 0.$$

Hence $\lambda_1 = \lambda_2 = \lambda_3 = \chi$, $2\chi = 0$, and $\lambda_0 = \chi$. The Fourier coefficient of an order-two character is real, and the trivial character contributes zero. The sum is precisely $E_{\langle 2 \rangle}(f)$. These cases exhaust all possibilities. $\square$

III.8.3 Counting the two types

Define

$$R_d = \frac{6^d - 2 \cdot 4^d + 2^d}{8}, \quad (\text{III.8.5})$$

$$F_d = \frac{2^d(2^d - 1)(2^d - 2)}{24}, \quad (\text{III.8.6})$$

$$T_d = F_d - R_d = \frac{8^d - 3 \cdot 6^d + 3 \cdot 4^d - 2^d}{24}. \quad (\text{III.8.7})$$

Here F_d is a combinatorial count and should not be confused with the function F .

Lemma III.8.2. *The numbers of ordinary parallelograms and binary affine planes in V_d are R_d and F_d , respectively. The number of non-parallelogram binary planes is T_d .*

Proof. For one binary coordinate, there are six ordered solutions of $a + b = c + e$: one with sum 0, four with sum 1, and one with sum 2. Thus there are 6^d ordered solutions in V_d . The repeated-vertex solutions are exactly those with $(a, b) = (c, e)$ or $(a, b) = (e, c)$, numbering $2 \cdot 4^d - 2^d$. Every remaining four-element parallelogram is counted eight times, corresponding to the orderings of its two diagonals and their interchange. This proves (III.8.5).

A binary affine plane has four possible origins and six ordered bases for its two-dimensional direction space. Choosing an origin and two independent directions in $\mathbb{F}_2^d$ gives $2^d(2^d - 1)(2^d - 2)$ choices, each plane occurring 24 times. This proves (III.8.6); subtraction gives (III.8.7). $\square$

d	$\binom{2^d}{4}$	R_d	T_d	F_d
2	1	1	0	1
3	70	12	2	14
4	1,820	100	40	140
5	35,960	720	520	1,240
6	635,376	4,816	5,600	10,416

The middle columns count all parallelograms, not just coordinate square faces. In particular, a three-dimensional Boolean cube contains twelve ordinary parallelograms, not six.

Theorem III.8.3 (Exact torsion-sensitive quartic expansion). *Under the hypotheses of Theorem III.7.3, put*

$$Q_d(f) = R_d E_4(f) + T_d E_{\langle 2 \rangle}(f).$$

Then

$$\mathcal{C}_d(\delta + f) = \delta^p + \delta^{p-4} Q_d(f) + \mathcal{R}_{\geq 5}, \quad (\text{III.8.8})$$

where

$$|\mathcal{R}_{\geq 5}| \leq \sum_{r=5}^p \binom{p}{r} \delta^{p-r} t^r. \quad (\text{III.8.9})$$

In particular,

$$\mathcal{D}_d(F) \leq F_d \delta^{p-4} t^4 + \sum_{r=5}^p \binom{p}{r} \delta^{p-r} t^r. \quad (\text{III.8.10})$$

If $|G|$ is odd, F_d in this upper bound can be replaced by R_d .

Proof. In (III.7.3), group the four-element subsets according to Lemma III.8.1 and use Lemma III.8.2. The remaining subsets have size at least five, so (III.6.8) gives (III.8.9). Finally, $E_{\langle 2 \rangle}(f) \leq E_4(f) \leq t^4$, with $E_{\langle 2 \rangle}(f) = 0$ in odd order. $\square$

Remark III.8.4 (What the torsion correction measures). For a group of exponent 2, the quartic coefficient is $F_d E_4(f)$. For an odd-order group it is $R_d E_4(f)$. In a general group, the additional contribution sees only the part of the Fourier spectrum lying on characters of order two. It is not a correction proportional to the total size of the two-torsion subgroup.

Remark III.8.5 (A sharp coefficient and a potentially non-sharp one). The universal leading coefficient F_d in (III.8.10) is sharp; a real order-two character attains it to fourth order (Proposition III.11.1). The coefficient R_d in the odd-order upper bound is the exact multiplier of $E_4(f)$, but its optimality as a multiplier of t^4 is a different question. We do not identify those two optimization problems.

[write] **6 October 2026; sources 06–10.** Theorem III.8.3 is source 03's Theorem III.2.6 (same R_d , T_d , and $E_{\langle 2 \rangle} = \mathcal{E}_{\langle 2 \rangle}$) and the quartic part of source 09's Theorem III.21.1 ($P_d = R_d$); three independent proofs, by Fourier kernels with binary minors here, by an XOR classification in source 03 and by lattice indices in source 09. The odd-order question in the last remark is source 03's Research question V.14.1 and this source's Question V.34.1; it is settled for $d = 3$ by source 04 (note there).

III.9 The complete three-cube expansion in odd order

Source 08, delivered Section 5.

In this section $d = 3$, $p = 8$, and G has odd order. Define

$$S_5(f) = \sum_{\chi \in \widehat{G}} |\widehat{f}(\chi)|^2 \widehat{f}(\chi)^2 \overline{\widehat{f}(2\chi)}. \quad (\text{III.9.1})$$

The sum is real: the summand at $-\chi$ is the complex conjugate of the summand at χ . Set

$$\begin{aligned} A_6(f) &= \Lambda_{V_3 \setminus \{000, 001\}}(f), \\ B_6(f) &= \Lambda_{V_3 \setminus \{000, 011\}}(f), \\ C_6(f) &= \Lambda_{V_3 \setminus \{000, 111\}}(f), \\ A_7(f) &= \Lambda_{V_3 \setminus \{000\}}(f). \end{aligned}$$

Binary strings denote vertices, and Λ_S is still averaged over all four cube parameters.

Lemma III.9.1 (Five-vertex terms). *For a real mean-zero f on an odd-order group,*

$$\sum_{\substack{S \subseteq V_3 \\ |S|=5}} \Lambda_S(f) = 8S_5(f).$$

Proof. There are $\binom{8}{5} = 56$ five-element subsets. Exactly 48 contain an ordinary parallelogram: there are 12 parallelograms, each with four choices of an additional vertex, and no five-element set contains two different parallelograms. To justify the last claim, two parallelograms sharing three vertices lie in the same affine plane; an affine two-plane contains at most four vertices of a Boolean cube. One way to see the latter fact is to choose two coordinate projections that are injective on its two-dimensional direction space; they map its binary vertices injectively into $\{0, 1\}^2$.

A parallelogram together with a fifth vertex has just its parallelogram relation as a rational dependence. More importantly, this remains the only dependence over $\widehat{G}$. Choose three affinely independent vertices of the parallelogram and the fifth vertex. Their 4×4 affine matrix has determinant ± 1 or ± 2 : reflect one chosen vertex to zero and use the binary 3×3 determinant bound from the proof of Lemma III.8.1. Multiplication by either determinant is invertible on an odd-order group. Subtracting a suitable multiple of the primitive parallelogram relation from any Fourier solution kills one parallelogram frequency, after which this invertible matrix kills all remaining frequencies. In particular, the fifth frequency is zero. Its factor $\widehat{f}(0)$ makes the average vanish.

The other eight five-element sets are the cube symmetries of

$$S_* = \{000, 110, 101, 011, 111\}.$$

Indeed, for each of the eight vertices w , delete the three neighbors of w . These give eight distinct sets with no parallelogram. To check the latter property for S_* , its unique rational dependence, in the displayed order, is

$$(-1, 1, 1, 1, -2), \tag{III.9.2}$$

which has full support. A parallelogram would give a relation with one zero coefficient. The count $56 - 48 = 8$ proves exhaustion.

For S_* , the three coordinate equations show directly, over any character group, that the frequencies at 110, 101, 011 coincide, say at χ . The frequencies at 111 and 000 are -2χ and $-\chi$, respectively. Thus

$$\Lambda_{S_*}(f) = \sum_{\chi} \widehat{f}(-\chi) \widehat{f}(\chi)^3 \widehat{f}(-2\chi) = S_5(f).$$

Cube reflections and coordinate permutations preserve the average, proving the asserted factor of eight. $\square$

Theorem III.9.2 (Exact three-cube polynomial). *Let G have odd order and $f : G \rightarrow \mathbb{R}$ have mean zero. For $\delta \geq 0$,*

$$\begin{aligned} \mathcal{C}_3(\delta + f) &= \delta^8 + 12\delta^4 E_4(f) + 8\delta^3 S_5(f) \\ &\quad + \delta^2(12A_6(f) + 12B_6(f) + 4C_6(f)) \\ &\quad + 8\delta A_7(f) + \mathcal{C}_3(f). \end{aligned} \tag{III.9.3}$$

Consequently, with $t = \|f\|_{U^3(G)}$,

$$0 \leq \mathcal{D}_3(\delta + f) \leq 12\delta^4 t^4 + 8\delta^3 t^5 + 28\delta^2 t^6 + 8\delta t^7 + t^8. \tag{III.9.4}$$

Proof. The constant and degree-four terms follow from Theorem III.8.3, since $R_3 = 12$ and there is no order-two Fourier contribution. Lemma III.9.1 gives the fifth-degree term. A six-element subset is determined by its missing pair of vertices. Pairs at Hamming distances 1, 2, 3 form three symmetry classes of sizes 12, 12, 4, respectively. These give the three degree-six averages displayed above. The eight seven-element subsets are all equivalent under cube reflections, and the eighth-degree term is the full cube of f . This proves the identity. Apply (III.6.8) to each average, and use $E_4(f) \leq t^4$. $\square$

[write] **6 October 2026; sources 11–14.** Source 13’s Theorem III.31.5 is this identity over $\mathbb{F}_p^n$, p odd, written with the orbit moments T_5, B_1, B_2, B_3, T_7 , which it bounds (including a spectral deficit); source 14’s Appendix A (Section III.38) certifies the same support counts.

The naive fourth-order estimate in dimension three has coefficients 70, 56, 28, 8, 1 in degrees 4, 5, 6, 7, 8. The all-group quartic analysis changes the first coefficient to 14. In odd order the exact support analysis improves these coefficients to 12, 8, 28, 8, 1. In particular, if $t \leq \delta$,

$$\mathcal{D}_3(\delta + f) \leq 57\delta^4 t^4. \tag{III.9.5}$$

This is a convenient uniform corollary, not an assertion that 57 is the best possible constant.

[write] **6 October 2026; sources 06–10.** Theorem III.9.2 contains source 04’s Lemma III.3.3 (the same support counts 12, 8, 28, 8, 1; this S_5 is 04’s B_5) and is the $d = 3$ case of source 09’s Theorem III.21.1 in odd order. Source 04’s Theorem III.3.6 replaces the leading $12\delta^4 t^4$ by the sharp $6\sqrt{2}\delta^4 t^4$.

III.9.1 A Fourier-sensitive remainder

Let $m(f) = \max_\chi |\widehat{f}(\chi)|$. Directly from (III.9.1),

$$|S_5(f)| \leq m(f)E_4(f). \tag{III.9.6}$$

Together with the cube inequality, we also have $|S_5(f)| \leq t^5$. Thus (III.9.3) gives the two-sided estimate

$$\left| \mathcal{D}_3(\delta + f) - 12\delta^4 E_4(f) \right| \leq 8\delta^3 m(f)E_4(f) + 28\delta^2 t^6 + 8\delta t^7 + t^8. \tag{III.9.7}$$

This is useful when the fourth Fourier moment is substantially smaller than t^4 . It does not turn U^2 uniformity alone into control of the full three-cube count: the remaining t^6, t^7, t^8 terms still matter.

The formula also identifies a concrete interaction invisible in a simple norm bound. The fifth-order contribution couples frequencies χ and 2χ . Its sign depends on their relative phases. Even for real weights, that sign need not be positive.

III.10 A counterexample to the quartic global lower bound

Source 08, delivered Section 6.

The positive fourth-order term might suggest the stronger lower estimate

$$\mathcal{D}_3(F) \stackrel{?}{\geq} 12\delta^4 E_4(F - \delta) \quad (|G| \text{ odd}, \quad 0 \leq F \leq 1). \quad (\text{III.10.1})$$

It is false. This is a natural strengthening posed and settled within this memorandum, not a claim to settle a previously identified published conjecture.

Theorem III.10.1 (Negative fifth-order correction). *On $G = \mathbb{Z}/11\mathbb{Z}$, put*

$$h(x) = \cos(2\pi x/11) - \cos(4\pi x/11).$$

For every fixed $0 < \delta < 1$, the weight $F = \delta + \eta h$ belongs to $[0, 1]$ and violates (III.10.1) for all sufficiently small positive η .

Proof. The nonzero Fourier coefficients of h are

$$\hat{h}(1) = \hat{h}(-1) = \frac{1}{2}, \quad \hat{h}(2) = \hat{h}(-2) = -\frac{1}{2}.$$

Thus $\mathbb{E}h = 0$, $E_4(h) = 1/4$, and (III.9.1) gives $S_5(h) = -1/16$: only $\chi = 1, -1$ contribute, each with value $-1/32$. Applying (III.9.3) to $f = \eta h$ yields

$$\mathcal{D}_3(\delta + \eta h) - 12\delta^4 E_4(\eta h) = -\frac{1}{2}\delta^3 \eta^5 + O_\delta(\eta^6), \quad (\text{III.10.2})$$

which is negative for sufficiently small positive η . Since $|h| \leq 2$, the condition $2\eta \leq \min(\delta, 1 - \delta)$ guarantees $0 \leq F \leq 1$.

There is also a completely explicit safe range. Use $\|h\|_{U^3} \leq 2$ in the terms of degree at least six in (III.9.3). Their absolute value, divided by $\frac{1}{2}\delta^3 \eta^5$, is at most

$$3584 \frac{\eta}{\delta} + 2048 \left(\frac{\eta}{\delta}\right)^2 + 512 \left(\frac{\eta}{\delta}\right)^3.$$

This is less than 1 when $0 < \eta/\delta \leq 10^{-4}$. Combining that condition with the pointwise range condition proves the claim without any numerical evaluation of a Gowers norm. $\square$

III.10.1 An exact finite certificate and a larger explicit example

For the same function h , a finite Fourier calculation gives

$$\mathcal{C}_3(\delta + \eta h) = \delta^8 + 3\delta^4 \eta^4 - \frac{1}{2}\delta^3 \eta^5 + \frac{3}{2}\delta^2 \eta^6 + \frac{1}{4}\eta^8. \quad (\text{III.10.3})$$

Here is a compact way to reproduce and check that calculation. The four Fourier constraints on the eight cube vertices have the following complete parametrization, with arithmetic modulo 11:

v	000	001	010	100	011	101	110	111
λ_v	a	b	c	e	$-a - b - c$	$-a - b - e$	$-a - c - e$	$2a + b + c + e$.

(III.10.4)

Every frequency must lie in $D = \{0, 1, -1, 2, -2\}$. A zero frequency contributes δ ; a frequency ± 1 contributes $\eta/2$; a frequency ± 2 contributes $-\eta/2$. It therefore suffices to check $5^4 = 625$ tuples $(a, b, c, e) \in D^4$. The resulting signed counts are

Degree r	0	1	2	3	4	5	6	7	8
Positive products	1	0	0	0	48	0	96	0	64
Negative products	0	0	0	0	0	16	0	0	0

Dividing the signed count in degree r by 2^r gives (III.10.3). Appendix III.15 supplies the short exhaustive certificate; the accompanying verification program also checks it against an independent eight-frequency enumeration. The general negative-sign conclusion in Theorem III.10.1 does not depend on this enumeration.

For an explicit example with comfortable margins from 0 and 1, take $\delta = 1/2$ and $\eta = 1/12$. Then $1/3 \leq F \leq 2/3$ and

$$\frac{\mathcal{D}_3(F)}{\delta^4 E_4(F - \delta)} = 12 - 2s + 6s^2 + s^4 \Big|_{s=1/6} = \frac{15337}{1296} < 12. \quad (\text{III.10.5})$$

The distinction between a local Taylor coefficient and a global lower-bound coefficient is therefore substantive, not just a possible technical caveat.

[write] **6 October 2026; sources 06–10.** Kept on record as Remark W1 in source 03’s section, where source 03 notes the same failure on $\mathbb{Z}/3\mathbb{Z}$; source 09’s Proposition III.22.2 proves it on primes $q \geq 17$. The polynomial (III.10.3) and the ratio (III.10.5) were recomputed independently at the second write (all 11^4 frequency choices, exact arithmetic).

III.11 Sharpness of the fourth-order exponent

Source 08, delivered Section 7.

III.11.1 An exact order-two model

Proposition III.11.1. *Suppose G has a nontrivial character $\chi : G \rightarrow \{1, -1\}$. For $p = 2^d$, $d \geq 2$,*

$$\mathcal{C}_d(\delta + \eta\chi) = \frac{(\delta + \eta)^p + (\delta - \eta)^p + 2(p-1)(\delta^2 - \eta^2)^{p/2}}{2p}. \quad (\text{III.11.1})$$

Moreover, $\|\eta\chi\|_{U^d} = |\eta|$, and for fixed $\delta > 0$,

$$\mathcal{D}_d(\delta + \eta\chi) = F_d \delta^{p-4} \eta^4 + O_{d,\delta}(\eta^6).$$

Thus the leading constant F_d in the all-group bound (III.8.10) is sharp.

Proof. The $d+1$ signs $\chi(x), \chi(h_1), \dots, \chi(h_d)$ are independent and uniform. If every increment sign is 1, all cube values have the same character sign; the two choices of $\chi(x)$ give the first two terms in the numerator. There are $p-1$ other increment-sign patterns. For each, exactly $p/2$ vertices have each sign, so either base sign gives the product $(\delta^2 - \eta^2)^{p/2}$. This proves the exact formula.

A character has Gowers norm 1 in every dimension at least two. The coefficient of η^4 in (III.11.1) is

$$\frac{2\binom{p}{4} + 2(p-1)\binom{p/2}{2}}{2p} = \frac{p(p-1)(p-2)}{24} = F_d.$$

The coefficient of η^2 cancels and all odd coefficients vanish. The remainder is of degree at least six, with the zero remainder understood when $p = 4$. $\square$

[write] **6 October 2026; sources 06–10.** (III.11.1) is source 09’s (III.22.1) (the two closed forms were expanded and compared at the second write for $d = 2, \dots, 5$). Its consequence, that the universal coefficient F_d is attained, is also source 03’s optimality of K_d over all groups and source 09’s Proposition III.22.1.

III.11.2 Odd cyclic groups and real cosine perturbations

To show that the exponent does not rely on two-torsion, let $N > p$ be an odd prime and put

$$h_N(x) = \cos(2\pi x/N), \quad x \in \mathbb{Z}/N\mathbb{Z}.$$

Then $\mathbb{E}h_N = 0$ and $E_4(h_N) = 1/8$. There is a positive constant c_d , depending only on d , such that

$$\|h_N\|_{U^d} = c_d \quad (N > p). \quad (\text{III.11.2})$$

Indeed, in the Fourier expansion of its full cube, every frequency is 1 or -1 . Every constraint is an integer of absolute value at most p ; since $N > p$, vanishing modulo N is equivalent to integer vanishing. The number of admissible frequency assignments is therefore independent of N . There is at least one assignment, namely the alternating cube assignment $\lambda_v = (-1)^{|v|}$, and every summand is positive. This proves (III.11.2) and $c_d > 0$.

For fixed $0 < \delta < 1$ and sufficiently small $\eta > 0$, the weight $F_N = \delta + \eta h_N$ takes values in $[0, 1]$. By Theorem III.8.3, uniformly for $N > p$,

$$\mathcal{D}_d(F_N) = \frac{R_d}{8} \delta^{p-4} \eta^4 + O_{d,\delta}(\eta^5), \quad \|F_N - \delta\|_{U^d}^p = c_d^p \eta^p. \quad (\text{III.11.3})$$

Since $R_d > 0$, no exponent larger than $4/p$ can replace $4/p$ in a uniform bound for these bounded real weights. The remainder could be improved in this particular example, but degree five already suffices for sharpness.

III.11.3 A rounding lemma in the Gowers norm

Weighted sharpness does not automatically imply sharpness for sets, so a separate argument is needed.

Lemma III.11.2 (Independent rounding). *Let $F : G \rightarrow [0, 1]$, $|G| = N$, and $d \geq 2$, $p = 2^d$. There exists a set $A \subseteq G$ such that*

$$\|\mathbf{1}_A - F\|_{U^d(G)} \leq \left(\frac{2\binom{p}{2}}{N} \right)^{1/p}. \quad (\text{III.11.4})$$

Proof. Choose the random variables $\mathbf{1}_A(y)$ independently, with success probabilities $F(y)$, and set $X(y) = \mathbf{1}_A(y) - F(y)$. They have mean zero under this random choice and satisfy $|X(y)| \leq 1$. For a fixed cube whose p sampled vertices are distinct, independence implies

$$\mathbb{E}_{\text{rand}} \prod_{v \in V_d} X(L_v) = 0.$$

For any pair of distinct Boolean vertices v, w , the difference $L_v - L_w$ is a linear form in the increments with at least one coefficient 1 or -1 . It is uniform in G , so $\Pr_{x,h}(L_v = L_w) = 1/N$. By the union bound, the probability that the sampled cube has any collision is at most $\binom{p}{2}/N$.

The product is bounded in absolute value by 1 in all cases. Interchanging the two finite averages therefore gives

$$0 \leq \mathbb{E}_{\text{rand}} \|X\|_{U^d}^p \leq \frac{\binom{p}{2}}{N}.$$

Markov's inequality shows that some realization has the moment at most twice this bound, which is (III.11.4). The random centering used here is pointwise in the probability space; it need not give $\mathbb{E}_{y \in G} X(y) = 0$ for an individual realization. $\square$

[write] **6 October 2026; sources 06–10.** Source 09's Lemma III.22.4 is the same random-rounding lemma with the constant $(\binom{p}{2}/|G|)^{1/p}$ (without the factor two of the Markov step); source 03 rounds on $(\mathbb{Z}/2\mathbb{Z})^n$ and $(\mathbb{Z}/3\mathbb{Z})^n$ with an exact density and source 04 on prime cyclic groups by variance. Four routes to sharpness for sets.

Lemma III.11.3 (Continuity of cube excess). *If $A \subseteq G$, $F : G \rightarrow [0, 1]$, $\delta = \mathbb{E}F$, $\delta_A = |A|/|G|$, and $q = \|\mathbf{1}_A - F\|_{U^d}$, then*

$$|\delta_A - \delta| \leq q, \quad (\text{III.11.5})$$

$$|\mathcal{D}_d(\mathbf{1}_A) - \mathcal{D}_d(F)| \leq 2pq, \quad (\text{III.11.6})$$

$$\| \|\mathbf{1}_A - \delta_A\|_{U^d} - \|F - \delta\|_{U^d} \| \leq 2q. \quad (\text{III.11.7})$$

Proof. The first assertion is (III.6.5) applied to $X = \mathbf{1}_A - F$. Telescope the difference between the two p -factor cube products. Each of the p terms has one factor X , while all other factors are F or $\mathbf{1}_A$, whose U^d norms are at most 1. Equation (III.6.4) gives $|\mathcal{C}_d(\mathbf{1}_A) - \mathcal{C}_d(F)| \leq pq$. Since $\delta, \delta_A \in [0, 1]$, $|\delta_A^p - \delta^p| \leq p|\delta_A - \delta| \leq pq$. This proves the excess estimate. Finally the difference of the centered functions is $X - (\delta_A - \delta)$, whose norm is at most $2q$ by Minkowski and the constant-function norm. The reverse triangle inequality proves the last claim. $\square$

Theorem III.11.4 (Sharpness for genuine sets in odd prime order). *Fix $d \geq 2$, $p = 2^d$, and $0 < \delta < 1$. There are primes $N_j \rightarrow \infty$ and sets $A_j \subseteq \mathbb{Z}/N_j\mathbb{Z}$ such that, with $\delta_j = |A_j|/N_j$ and $\alpha_j = \|\mathbf{1}_{A_j} - \delta_j\|_{U^d}^p$,*

$$\delta_j \rightarrow \delta, \quad \alpha_j \rightarrow 0, \quad \frac{\mathcal{D}_d(\mathbf{1}_{A_j})}{\alpha_j^{4/p}} \rightarrow \frac{R_d \delta^{p-4}}{8c_d^4} > 0.$$

For every $\beta > 4/p$, $\mathcal{D}_d(\mathbf{1}_{A_j})/\alpha_j^\beta \rightarrow \infty$. Thus a bound uniform over such sets, even with density restricted to any fixed neighborhood of δ , cannot have exponent larger than $4/p$.

Proof. Choose $\eta_j = c/j$ with $0 < c < \min(\delta, 1 - \delta)$. Choose a prime $N_j > \max(p, j^{8p})$ and set $F_j = \delta + \eta_j h_{N_j}$. Apply Lemma III.11.2 to obtain A_j with rounding error $q_j = O_d(j^{-8}) = o(\eta_j^4)$. By Lemma III.11.3 and (III.11.3),

$$\begin{aligned} \delta_j &= \delta + o(1), \\ \|\mathbf{1}_{A_j} - \delta_j\|_{U^d} &= c_d \eta_j + o(\eta_j), \\ \mathcal{D}_d(\mathbf{1}_{A_j}) &= \frac{R_d}{8} \delta^{p-4} \eta_j^4 + o(\eta_j^4). \end{aligned}$$

Divide the last relation by the fourth power of the norm in the second. For $\beta > 4/p$, the additional factor $\alpha_j^{4/p-\beta}$ tends to infinity. $\square$

Remark III.11.5 (The density statement is intentionally asymptotic). The sets in this theorem have densities tending to the prescribed δ ; we do not assert exact density δ in every prime cyclic group. For instance, density $1/2$ is impossible in an odd-order group. This does not weaken sharpness for a uniform bound such as (III.7.5).

III.12 Lower bounds and Fourier stability

Source 08, delivered Section 8.

The preceding upper estimates can be paired with a strengthened form of the basic lower cube count. This section records consequences of norm monotonicity and Parseval, rather than attributing a new general norm theorem to the present calculation.

Theorem III.12.1 (Fourier-moment lower bound). *For a real function $F = \delta + f$ on a finite abelian group, where $\delta \geq 0$ and $\mathbb{E}f = 0$, and for $d \geq 2$, $p = 2^d$,*

$$\mathcal{C}_d(F) \geq (\delta^4 + E_4(f))^{p/4}. \quad (\text{III.12.1})$$

In particular,

$$\mathcal{D}_d(F) \geq \frac{p}{4} \delta^{p-4} E_4(f). \quad (\text{III.12.2})$$

For $\delta > 0$, if $D = \mathcal{D}_d(F)$, then

$$E_4(f) \leq (\delta^p + D)^{4/p} - \delta^4 \leq \frac{4}{p} \delta^{4-p} D. \quad (\text{III.12.3})$$

Equality in $\mathcal{C}_d(F) \geq \delta^p$ holds if and only if F is constant.

Proof. The constant frequency of F is δ , while all other Fourier coefficients are those of f . Thus $\|F\|_{U^2}^4 = \delta^4 + E_4(f)$. Now $\|F\|_{U^d} \geq \|F\|_{U^2}$, giving (III.12.1). The binomial theorem, since $p/4$ is a positive integer, implies (III.12.2). Rearrange (III.12.1) to obtain the first inequality in (III.12.3). The concavity of $u \mapsto u^{4/p}$ gives the second inequality by the tangent line at $u = \delta^p$. If the cube count equals δ^p , then (III.12.1) forces $E_4(f) = 0$. Every Fourier coefficient of f is zero, so $f = 0$. The converse is immediate. $\square$

As a direct stability statement, every nontrivial Fourier coefficient obeys

$$\max_{\chi \neq 0} |\widehat{F}(\chi)| \leq \left(\frac{4}{p} \delta^{4-p} \mathcal{D}_d(F) \right)^{1/4} \quad (\delta > 0).$$

This is stability in Fourier space. It is not closeness in L^2 to a constant: for an indicator, the squared L^2 distance is exactly $\delta(1 - \delta)$, even along the near-minimal sequences in Theorem III.11.4. The displayed inequality itself controls U^2 , not the balanced U^d norm; a stronger norm-geometric conclusion would require an additional argument.

Corollary III.12.2 (A finite-size correction). *Suppose $|G| = N > 1$ and write $\sigma^2 = \mathbb{E}|F - \delta|^2$. Then*

$$\mathcal{C}_d(F) \geq \left(\delta^4 + \frac{\sigma^4}{N-1} \right)^{p/4}. \quad (\text{III.12.4})$$

For $A \subseteq G$ of density δ ,

$$\mathcal{C}_d(\mathbf{1}_A) \geq \left(\delta^4 + \frac{\delta^2(1-\delta)^2}{N-1} \right)^{p/4}. \quad (\text{III.12.5})$$

Proof. Parseval and $\widehat{f}(0) = 0$ give $\sum_{\chi \neq 0} |\widehat{f}(\chi)|^2 = \sigma^2$. Cauchy–Schwarz over the $N - 1$ nontrivial characters yields $E_4(f) \geq \sigma^4/(N - 1)$. Insert this into (III.12.1). For an indicator, $\sigma^2 = \delta(1 - \delta)$. $\square$

The lower coefficient $p/4$ in (III.12.2) and the local coefficient R_d in odd order are different quantities. Theorem III.10.1 shows explicitly why one cannot simply replace the former with the latter. Determining the best global lower coefficient remains a separate optimization problem.

III.13 Different weights at different vertices

Source 08, delivered Section 9.

The centering argument is not limited to counting cubes in a single set. Suppose each vertex has a real weight

$$F_v = \delta_v + f_v, \quad \delta_v = \mathbb{E}F_v, \quad \mathbb{E}f_v = 0, \quad t_v = \|f_v\|_{U^d}.$$

The means may be different and may be negative.

Theorem III.13.1 (Inhomogeneous fourth-order estimate). *With the notation above,*

$$\left| \Lambda_d((F_v)) - \prod_{v \in V_d} \delta_v \right| \leq \sum_{\substack{S \subseteq V_d \\ |S| \geq 4}} \prod_{v \notin S} |\delta_v| \prod_{v \in S} t_v. \quad (\text{III.13.1})$$

All terms with one, two, or three centered factors vanish identically.

Proof. Expand the product by the subset of vertices at which the centered factor is chosen. Lemma III.7.1 makes every term of sizes one to three a product of zero means. Pad the remaining selected inputs with constant 1 in (III.6.4) and take absolute values. $\square$

The quartic classification is also valid with different real centered inputs. If $S = \{a, b, c, e\}$ is an ordinary parallelogram with $a + b = c + e$, the associated mixed fourth-order average is

$$\sum_{\chi} \widehat{f}_a(\chi) \widehat{f}_b(\chi) \overline{\widehat{f}_c(\chi)} \overline{\widehat{f}_e(\chi)}. \quad (\text{III.13.2})$$

For a non-parallelogram binary plane it is

$$\sum_{\substack{\chi \neq 0 \\ 2\chi = 0}} \prod_{v \in S} \widehat{f}_v(\chi). \quad (\text{III.13.3})$$

For every other four-set it vanishes. These statements follow from the same character equations as Lemma III.8.1. Hölder’s inequality bounds the absolute value of either expression by

$$\prod_{v \in S} \|f_v\|_{U^2}.$$

Thus one can improve (III.13.1) further by keeping the quartic Fourier expressions and using U^d norms only from degree five onward. Unlike the common-function quartic term, a mixed quartic expression need not be nonnegative. Accordingly, no lower bound by $\prod_v \delta_v$ is asserted in this inhomogeneous setting.

This version applies, for example, when each cube vertex must belong to a different set A_v . It isolates the contribution of the densities and shows that the first possible deviation involves four of the balanced vertex conditions, rather than just one.

III.14 Which centered terms can survive?

Source 08, delivered Section 10.

The four- and five-vertex calculations are instances of a general finite linear-algebra principle. It is useful both for organizing higher coefficients and for designing checkable certificates.

III.14.1 An exact character-kernel criterion

Let $a_1, \dots, a_r \in \mathbb{Z}^m$ be coefficient columns and let $L_j(z) = a_j \cdot z$ for $z \in G^m$. Let M be the matrix with these columns. For independently chosen functions $g_j : G \rightarrow \mathbb{C}$,

$$\mathbb{E}_{z \in G^m} \prod_{j=1}^r g_j(L_j(z)) = \sum_{\lambda \in \ker_{\widehat{G}} M} \prod_{j=1}^r \widehat{g}_j(\lambda_j), \quad (\text{III.14.1})$$

where $\ker_{\widehat{G}} M$ means the solutions of $M\lambda = 0$ in $\widehat{G}^r$. A vector has *full support* if none of its coordinates is zero.

Proposition III.14.1 (Universal vanishing criterion). *The following conditions are equivalent:*

- (i) *The kernel $\ker_{\widehat{G}} M$ has no full-support vector.*
- (ii) *The average in (III.14.1) is zero for every choice of complex mean-zero functions $g_1, \dots, g_r$.*
- (iii) *The average is zero for every common real mean-zero function $g_1 = \dots = g_r = f$.*

Proof. If no kernel vector has full support, every summand in (III.14.1) contains a factor $\widehat{g}_j(0) = 0$. Thus (i) implies (ii), and (ii) implies (iii).

Conversely, suppose $\lambda \in \ker_{\widehat{G}} M$ has full support. For the independently chosen complex functions one may simply take $g_j = e_{\lambda_j}$; their product along the forms is identically 1. To obtain a common real witness as well, take

$$f = \sum_{j=1}^r (e_{\lambda_j} + e_{-\lambda_j}).$$

This function is real and has mean zero. Its Fourier coefficients are all nonnegative real numbers, and every frequency λ_j has a strictly positive coefficient. Therefore all summands in (III.14.1) are nonnegative and the summand indexed by λ is positive. This contradicts (iii), proving (iii) implies (i). Rescaling gives bounded witness functions if desired. $\square$

For a cube subset S , use the matrix with columns $(1, v)$ for $v \in S$. Let $A_{d,r}(G)$ be the number of r -element subsets whose character kernel has a full-support vector. We then have the systematic improvement

$$|\mathcal{C}_d(\delta + f) - \delta^p| \leq \sum_{r=1}^p A_{d,r}(G) \delta^{p-r} t^r \quad (\delta \geq 0, \mathbb{E}f = 0). \quad (\text{III.14.2})$$

In particular, $A_{d,1} = A_{d,2} = A_{d,3} = 0$. For a nontrivial odd-order group $A_{d,4} = R_d$, while a nontrivial elementary abelian two-group has $A_{d,4} = F_d$. The coefficients $A_{d,r}$ are exact counts of potentially nonzero supports, but the resulting norm bound is not automatically optimal: the same function need not simultaneously saturate every surviving Gowers–Cauchy–Schwarz inequality.

III.14.2 The large-prime criterion and exceptional characteristics

The column-dependence structure of an integer matrix is often called its *column matroid*. Here only a simple part of that language is needed: a column is a *coloop* if it does not lie in the span of the other columns.

Proposition III.14.2. *Fix an integer matrix M with r columns. There is a finite set of primes $\mathcal{B}_M$ such that, for every prime $q \notin \mathcal{B}_M$ with $q > r$,*

$$\ker_{\mathbb{F}_q} M \text{ has a full-support vector} \iff M \text{ has no coloop over } \mathbb{Q}.$$

Proof. For M and for each matrix obtained by deleting one column, choose a nonzero minor of size equal to its rational rank; for a rank-zero matrix, use the empty minor 1. Let $\mathcal{B}_M$ contain the prime divisors of these finitely many nonzero integers. Outside $\mathcal{B}_M$, all these ranks are preserved modulo q , since the chosen minors remain nonzero and larger minors vanish identically over the integers. Consequently, the coloop condition agrees over $\mathbb{Q}$ and $\mathbb{F}_q$.

If the j th column is a coloop, every relation has $\lambda_j = 0$, so no full-support vector exists. Conversely, suppose there is no coloop and let $K = \ker_{\mathbb{F}_q} M$ have dimension k . For every j , the coordinate map $\lambda \mapsto \lambda_j$ is a nonzero linear functional on K . Its zero set has size q^{k-1} . The union of the r coordinate zero sets has at most $rq^{k-1} < q^k$ elements. Some vector in K avoids all of them. $\square$

The restriction $q > r$ is deliberate. Even when every coordinate functional is nonzero, finitely many coordinate hyperplanes can cover a vector space over a small field. Likewise, a rank computation over $\mathbb{Q}$ alone cannot detect the order-two quartic terms found in Section III.8. For a general finite abelian group, one must solve the character-group constraints themselves; integer normal forms and prime-primary decomposition provide one way to organize that computation.

This framework is not a substitute for true-complexity theory. It decides *identical vanishing caused by centering*. True-complexity results ask which uniformity norms control the remaining nonzero configuration averages [29, 14]. The two analyses can be used consecutively: first remove terms that vanish exactly, then apply the strongest available norm estimate to the terms that remain.

III.15 A short exact certificate for the counterexample

Source 08, delivered Appendix A.

The code below implements (III.10.4) directly. Every loop variable and every comparison is an integer. Because that parametrization is a bijection between four free frequencies and the Fourier kernel, the enumeration has no omitted solutions and no multiplicities to correct. The returned coefficients are exact rational numbers.

```
from collections import Counter
from fractions import Fraction
from itertools import product

D = {0, 1, 10, 2, 9} # residues 0, +1, -1, +2, -2 mod 11
sign = {1: 1, 10: 1, 2: -1, 9: -1}
counts = Counter()
for a, b, c, e in product(D, repeat=4):
```

```

frequencies = [
    a, b, c, e,
    (-a-b-c) % 11, (-a-b-e) % 11,
    (-a-c-e) % 11, (2*a+b+c+e) % 11,
]
if not all(k in D for k in frequencies):
    continue
degree = sum(k != 0 for k in frequencies)
product_sign = 1
for k in frequencies:
    if k != 0:
        product_sign *= sign[k]
counts[degree, product_sign] += 1

coefficients = [
    Fraction(counts[r, 1] - counts[r, -1], 2**r)
    for r in range(9)
]
assert coefficients == [
    1, 0, 0, 0, 3, Fraction(-1, 2), Fraction(3, 2), 0,
    Fraction(1, 4),
]

```

The coefficient of degree r multiplies $\delta^{8-r}\eta^r$. This certificate is elementary executable arithmetic, not a Lean kernel certificate. The full verification program includes additional independent checks and records their output in `verification_results.json`.

III.16 Normalized cube forms and the analytic input

Source 09, delivered Section 2.

All groups in this paper are finite abelian groups written additively; the trivial group is allowed. Averages are normalized. The character group $\hat{G}$ is written multiplicatively, with identity character 1. The normalized Fourier transform is

$$a_\chi = \hat{f}(\chi) = \mathbb{E}_{x \in G} f(x) \overline{\chi(x)}, \quad f(x) = \sum_{\chi \in \hat{G}} a_\chi \chi(x). \quad (\text{III.16.1})$$

For real f , $a_{\chi^{-1}} = \overline{a_\chi}$, and $\mathbb{E}f = 0$ means $a_1 = 0$.

Let $V_d = \{0, 1\}^d$, and put

$$L_v(x, h) = x + \sum_{i=1}^d v_i h_i \quad (x \in G, h \in G^d, v \in V_d).$$

For real F , define

$$\mathcal{C}_d(F) = \mathbb{E}_{x, h} \prod_{v \in V_d} F(L_v(x, h)). \quad (\text{III.16.2})$$

For complex functions the Gowers form has complex conjugation at odd vertices:

$$\Lambda_d((g_v)) = \mathbb{E}_{x, h} \prod_{v \in V_d} \mathcal{C}^{|v|} g_v(L_v(x, h)), \quad \|g\|_{U^d}^{2^d} = \Lambda_d((g)_v),$$

where $\mathcal{C}$ denotes conjugation. For real functions the conjugations disappear, so (III.16.2) is $\|F\|_{U^d}^{2^d}$. For $d \geq 2$ this is a norm power. The case $d = 1$ is the seminorm $|\mathbb{E}F|$ and will not be used for the main results.

Lemma III.16.1 (Gowers–Cauchy–Schwarz). *For arbitrary complex functions $g_v : G \rightarrow \mathbb{C}$,*

$$|\Lambda_d((g_v))| \leq \prod_{v \in V_d} \|g_v\|_{U^d}.$$

Consequently, if f is real and $S \subseteq V_d$, then

$$\left| \mathbb{E}_{x,h} \prod_{v \in S} f(L_v(x, h)) \right| \leq \|f\|_{U^d}^{|S|}. \quad (\text{III.16.3})$$

The first assertion is the normalized form of Gowers’ Lemma 3.8 [1]. Its iterated Cauchy–Schwarz proof works over every finite abelian group: at each coordinate, separate the two opposite faces, apply Cauchy–Schwarz, and duplicate the averaging variable. After all coordinates have been duplicated, each vertex function has its own full cube average. If one starts with the paper’s bounded version, scaling each of the finitely many functions proves the unrestricted version. The second assertion follows by setting $g_v = f$ on S and $g_v = 1$ elsewhere.

For completeness, the auxiliary norm facts needed below follow directly from this input. Taking one vertex function to be g and all others one gives $|\mathbb{E}g| \leq \|g\|_{U^d}$. Taking g on one face and one on the opposite face gives $\|g\|_{U^{d-1}} \leq \|g\|_{U^d}$ for $d \geq 2$, interpreting U^1 as above. Expanding the cube form of $g + h$ and applying Lemma III.16.1 termwise gives the triangle inequality. Finally, Fourier orthogonality gives

$$\|f\|_{U^2}^4 = \sum_{\chi \in \widehat{G}} |a_\chi|^4. \quad (\text{III.16.4})$$

These observations also prove $\mathcal{C}_d(F) \geq (\mathbb{E}F)^m$ for real F with nonnegative mean.

Definition III.16.2 (Subset moments and binomial tails). For $S \subseteq V_d$, put

$$T_S(f) = \mathbb{E}_{x,h} \prod_{v \in S} f(L_v(x, h)), \quad T_\emptyset(f) = 1.$$

For $0 \leq r \leq m$ and $\delta, \varepsilon \geq 0$, put

$$B_{m,r}(\delta, \varepsilon) = \sum_{s=r}^m \binom{m}{s} \delta^{m-s} \varepsilon^s.$$

The elementary but useful estimate

$$B_{m,r}(\delta, \varepsilon) \leq \binom{m}{r} \varepsilon^r (\delta + \varepsilon)^{m-r} \quad (\text{III.16.5})$$

follows by comparing the coefficient of $\delta^{m-s} \varepsilon^s$ on both sides: $\binom{m}{r} \binom{m-r}{s-r} = \binom{m}{s} \binom{s}{r} \geq \binom{m}{s}$. We use the usual convention that a zeroth power is one, including at zero. No division by the density is needed in the main estimates.

III.17 Balancing removes the first three orders

Source 09, delivered Section 3.

Lemma III.17.1 (Three-vertex uniformity). *Let $S \subseteq V_d$ with $1 \leq |S| \leq 3$. If (x, h) is uniform in G^{d+1} , then $(L_v(x, h))_{v \in S}$ is uniform in G^S . In particular, $T_S(f) = 0$ whenever $\mathbb{E}f = 0$.*

Proof. Flipping an abstract cube coordinate is an invertible change of variables: replace h_i by $-h_i$ and x by $x + h_i$ for every flipped coordinate, with all replacements understood simultaneously. Thus we may move one selected vertex to 0.

For two selected vertices, the second is a nonzero binary vector. One of its coordinates is one, so the two values are controlled by x and an independent group variable with coefficient one.

For three selected vertices, write the other two as distinct nonzero binary vectors u, v . After interchanging them if necessary, there is a coordinate i with $(u_i, v_i) = (1, 0)$. There is a coordinate j with $v_j = 1$. The two rows $(1, 0)$ and $(u_j, 1)$ have determinant one. Together with the base variable, these coordinates give an integer minor of determinant ± 1 for the three value maps. Fixing all other variables, the remaining map is an automorphism of G^3 followed by a translation. Hence its output is uniform. The one-vertex case is immediate. Independence now gives $T_S(f) = (\mathbb{E}f)^{|S|} = 0$. $\square$

The determinant-one argument is important. Rational linear independence alone would not imply this statement on arbitrary finite groups.

Theorem III.17.2 (Fourth-order replacement for Lemma 3.10). *Let $d \geq 2$, $m = 2^d$, $f : G \rightarrow \mathbb{R}$ have mean zero, and $\delta \geq 0$. Put $\varepsilon = \|f\|_{U^d}$. Then*

$$0 \leq \mathcal{C}_d(\delta + f) - \delta^m \leq B_{m,4}(\delta, \varepsilon) \leq \binom{m}{4} \varepsilon^4 (\delta + \varepsilon)^{m-4}. \quad (\text{III.17.1})$$

In particular, for $A \subseteq G$ of density δ and $\|\mathbf{1}_A - \delta\|_{U^d}^m \leq \alpha$,

$$\#\{(x, h) : L_v(x, h) \in A \ \forall v\} \leq |G|^{d+1} \left[\delta^m + \sum_{s=4}^m \binom{m}{s} \delta^{m-s} \alpha^{s/m} \right]. \quad (\text{III.17.2})$$

Proof. Multilinearity gives the exact identity

$$\mathcal{C}_d(\delta + f) = \sum_{S \subseteq V_d} \delta^{m-|S|} T_S(f). \quad (\text{III.17.3})$$

The empty subset contributes δ^m . Lemma III.17.1 removes all subsets of sizes one through three. Bound the others by (III.16.3) and count subsets of each size. The final inequality is (III.16.5); the lower bound was proved after Lemma III.16.1. For an indicator, the normalized count becomes the integer count after multiplication by $|G|^{d+1}$, and the polynomial upper bound is increasing in $\varepsilon \geq 0$. $\square$

[write] **6 October 2026; sources 06–10.** Theorem III.17.2 is also proved by sources 03, 04 and 08 (note after source 08's Theorem III.7.3); (III.17.2) is source 08's proposed formal target (V.32.1).

For $d = 2$ there is only one surviving subset. Fourier orthogonality therefore gives the exact identity

$$\mathcal{C}_2(\delta + f) = \delta^4 + \|f\|_{U^2}^4. \quad (\text{III.17.4})$$

For larger d , the surviving terms are not arbitrary. Most four- and five-vertex subsets still contribute zero. Their complete classification is the main task of the next sections.

III.18 Integer value lattices and Fourier constraints

Source 09, delivered Section 4.

III.18.1 A general exact formula

Order a subset $S = \{v_1, \dots, v_s\} \subseteq V_d$ and form the $(d+1) \times s$ integer matrix

$$M_S = \begin{pmatrix} 1 & \cdots & 1 \\ v_1 & \cdots & v_s \end{pmatrix}. \quad (\text{III.18.1})$$

Let $\mathcal{L}_S$ be the integer lattice generated by its rows in $\mathbb{Z}^s$. Define the character constraint group

$$K_G(S) = \left\{ (\chi_j) \in \widehat{G}^s : \prod_j \chi_j = 1, \quad \prod_j \chi_j^{(v_j)_i} = 1 \ (1 \leq i \leq d) \right\}. \quad (\text{III.18.2})$$

Proposition III.18.1 (Fourier–lattice identity). *For functions $f_j : G \rightarrow \mathbb{C}$ with Fourier coefficients $a_{j,\chi}$,*

$$\mathbb{E}_{x,h} \prod_{j=1}^s f_j(L_{v_j}(x, h)) = \sum_{(\chi_j) \in K_G(S)} \prod_{j=1}^s a_{j,\chi_j}. \quad (\text{III.18.3})$$

If all f_j have mean zero, only tuples with every $\chi_j \neq 1$ can contribute.

Proof. Insert Fourier inversion for every factor. The average over x vanishes unless $\prod_j \chi_j = 1$. Independently, the average over h_i vanishes unless $\prod_j \chi_j^{(v_j)_i} = 1$. Every surviving character average is one. All sums are finite, so no interchange issue arises. $\square$

If $\mathcal{L}_S = \mathbb{Z}^s$, then the constraints force all $\chi_j = 1$. More generally, when M_S has full column rank, the quotient $\mathbb{Z}^s/\mathcal{L}_S$ is finite; its invariant factors specify the possible torsion characters. For example, an index-two quotient permits only one order-two parameter. An index-three quotient permits only one order-three parameter. The support of the corresponding relation still matters: a forced trivial character at even one vertex kills the moment.

When the rank is $s-1$, there can instead be a free character parameter coming from a real affine relation. If the row lattice is saturated in its rational span, that relation describes the entire constraint group, without additional torsion components. These are the cases that give ordinary fourth moments and the quintic statistic below.

III.18.2 Small determinant bounds

Lemma III.18.2. *Every 2×2 binary matrix has determinant of absolute value at most one; every 3×3 binary matrix has determinant of absolute value at most two; every 4×4 binary matrix has determinant of absolute value at most three.*

Proof. The first claim follows from the determinant formula. For an $n \times n$ binary matrix B , introduce the $(n+1) \times (n+1)$ sign matrix

$$H = \begin{pmatrix} 1 & \mathbf{1}^T \\ \mathbf{1} & J - 2B \end{pmatrix}.$$

Subtracting its first row from each remaining row shows $\det H = (-2)^n \det B$. Hadamard's determinant inequality gives

$$|\det B| \leq \frac{(n+1)^{(n+1)/2}}{2^n}.$$

For $n = 3$ this is two; for $n = 4$ it is less than four, so the integral determinant is at most three in absolute value. $\square$

Any minor of M_S which includes the all-ones row can be reduced, after moving one vertex to zero by cube-coordinate flips, to a binary minor one size smaller. This permits Lemma III.18.2 to control both primitive affine relations and possible lattice indices.

Lemma III.18.3 (A plane contains at most four cube vertices). *An affine subspace of real dimension r contains at most 2^r vertices of a Boolean cube. In particular, five distinct cube vertices have augmented rank at least four.*

Proof. Choose r coordinate functionals whose restrictions are independent on the direction space. Projection onto these coordinates is injective on the affine subspace. A cube vertex projects to an element of $\{0, 1\}^r$, of which there are 2^r . $\square$

III.19 The complete four-vertex classification

Source 09, delivered Section 5.

Vectors denoted by u, v, w, z in a normal form below are binary vectors with pairwise disjoint, nonempty supports. Coordinates unused by any of them are permitted. Cube symmetries mean coordinate permutations and coordinate flips, not arbitrary affine transformations of the cube.

Theorem III.19.1 (Four selected vertices). *For real mean-zero f , a four-vertex moment can be nonzero only in the following two types, up to cube symmetries:*

Type	Normal form	Moment
Parallelogram	$\{0, u, v, u + v\}$	$\sum_{\chi} a_{\chi} ^4$
Binary torsion	$\{0, u + v, u + w, v + w\}$	$\sum_{\chi^2=1} a_{\chi} ^4$

All other four-vertex moments vanish over every finite abelian group.

Proof. Any three selected vertices have augmented rank three and a primitive minor, by Lemma III.17.1. Thus the four-vertex rank is three or four.

Suppose first that it is three. A primitive relation has no zero coefficient, since any three vertices are independent. Choose three independent rows including the all-ones row. The cofactors giving its relation have absolute value at most one: after a cube flip, they are 2×2 binary determinants. Hence the relation has four coefficients ± 1 . Their sum is zero, so two have each sign. Moving one vertex to zero reduces the relation to $u + v = w$, and binarity forces u and v to have disjoint supports. Both are nonempty. The value relation is an ordinary parallelogram relation, and the primitive three-vertex minor excludes additional torsion. Formula (III.18.3) gives $\sum_{\chi} a_{\chi}^2 a_{\chi^{-1}}^2 = \sum_{\chi} |a_{\chi}|^4$.

Suppose next that the rank is four. Some full minor including the all-ones row is nonzero, and it has absolute value at most two. Therefore the lattice index is one or two. Index one makes the moment zero. At index two the nonzero relation modulo two is unique. Its support has even size because of the all-ones row; it cannot have size two because the vertices are distinct. It therefore uses all four vertices.

Move one vertex to zero. In each coordinate, the remaining three bits have even sum. The only patterns are 000, 110, 101, 011. Full rank requires all three nonzero patterns to occur. Grouping equal patterns produces precisely $\{0, u + v, u + w, v + w\}$ with three disjoint nonempty blocks. Its quotient has order two, and its character constraints are (χ, χ, χ, χ) with $\chi^2 = 1$. Such a Fourier coefficient is real, so its fourth power is $|a_\chi|^4$. $\square$

Proposition III.19.2 (Counting the four-vertex types). *The numbers of parallelogram and binary-torsion subsets of V_d are respectively*

$$P_d = \frac{2^d}{8}(3^d - 2 \cdot 2^d + 1) = \frac{6^d - 2 \cdot 4^d + 2^d}{8}, \quad (\text{III.19.1})$$

$$T_d = \frac{2^d}{24}(4^d - 3 \cdot 3^d + 3 \cdot 2^d - 1) = \frac{8^d - 3 \cdot 6^d + 3 \cdot 4^d - 2^d}{24}. \quad (\text{III.19.2})$$

Proof. For a parallelogram, choose a base vertex in 2^d ways and assign each coordinate to the unused block, the support of u , or the support of v , requiring both active blocks to occur. Inclusion–exclusion gives $3^d - 2 \cdot 2^d + 1$ assignments. Each resulting subset is counted for four choices of base and two orders of the active blocks, so divide by eight.

For the second type, choose a base and assign coordinates to an unused block or one of three nonempty active blocks. There are $4^d - 3 \cdot 3^d + 3 \cdot 2^d - 1$ assignments. Each subset has four base choices and $3!$ orders of the blocks, so divide by twenty-four. $\square$

In particular, the fourth-order coefficient is never negative. Its arithmetic dependence is entirely through the order-two part of the character group. There are no other fourth-order corrections, even on composite groups with large exponents.

[write] **6 October 2026; sources 06–10.** Theorem III.19.1 and Proposition III.19.2 are source 03’s Lemma III.2.3 with its counts and source 08’s Lemmas III.8.1 and III.8.2: P_d is 03’s and 08’s R_d , T_d is theirs, and $P_d + T_d = A_d$ is 08’s F_d and 03’s K_d (checked for $d = 2, \dots, 7$ at the second write). Three independent proofs.

III.20 The complete five-vertex classification

Source 09, delivered Section 6.

Theorem III.20.1 (Five selected vertices). *For real mean-zero f , every five-vertex moment vanishes except possibly for the following two types, up to cube symmetries:*

$$\text{star:} \quad \{0, u, v, w, u + v + w\}, \quad (\text{III.20.1})$$

$$\text{ternary torsion:} \quad \{0, u + v + w, u + v + z, u + w + z, v + w + z\}. \quad (\text{III.20.2})$$

The star moment is

$$\sum_{\chi \in \widehat{G}} |a_\chi|^2 a_\chi^2 \overline{a_{\chi^2}}, \quad (\text{III.20.3})$$

and the ternary-torsion moment is

$$\sum_{\chi^3=1} a_\chi^4 \overline{a_\chi}. \quad (\text{III.20.4})$$

Proof. By Lemma III.18.3, the augmented rank is four or five.

Rank four, with a zero in the unique relation. The relation must involve exactly four vertices, and those four form a parallelogram. Put them in the form $0, u, v, u + v$, with u, v supported on disjoint nonempty coordinate blocks U, V . Let w be the fifth vertex. It lies outside their real affine plane. Therefore at least one of the following holds: w is nonzero outside $U \cup V$; its bits vary within U ; or its bits vary within V .

In the first case, a coordinate outside $U \cup V$, together with one coordinate in each of U, V , gives a determinant-one minor for $0, u, v, w$. In the second case, take two coordinates in U where the bits of w differ and one in V ; their determinant is ± 1 . The third case is symmetric. Thus the four values at $0, u, v, w$ are jointly uniform over every group. The fourth parallelogram value is determined by the first three, while the value at w is independent. Its mean-zero factor kills the entire moment. This argument rules out hidden torsion, rather than merely using rational independence.

Rank four, with all five relation coefficients nonzero. Choose four independent rows including the all-ones row. Every cofactor has absolute value at most two, by Lemma III.18.2. The primitive relation therefore has coefficients of absolute values one or two, with sum zero. Up to an overall sign and permutation, the only possibilities are

$$(1, 1, 1, -1, -2), \quad (2, 2, -2, -1, -1).$$

Indeed, the number of odd coefficients is even; a primitive vector cannot have all coefficients even. Enumerating one or three even coefficients gives these two possibilities.

The second is impossible for distinct cube vertices. In any coordinate, the zero-sum condition forces the two vertices with coefficient -1 to have equal bits, since all other contributions are even. Those two vertices would therefore coincide. For the first relation, move the coefficient -2 vertex to zero. If u, v, w are the coefficient-1 vertices and a is the coefficient -1 vertex, the relation reads $u + v + w = a$. Binariness forces the supports of u, v, w to be disjoint; distinctness makes them nonempty. This gives (III.20.1).

There is no extra torsion in this case. The cofactor vector is an integer multiple of the primitive relation. Since one coefficient has absolute value two and all cofactors have absolute value at most two, that multiple has absolute value one. Some cofactor is consequently ± 1 , so the row lattice is saturated. In the order $(0, u, v, w, u + v + w)$, a primitive relation is $(-2, 1, 1, 1, -1)$. The character tuple is therefore $(\chi^{-2}, \chi, \chi, \chi, \chi^{-1})$. Its Fourier product is $a_\chi^3 \overline{a_\chi} \overline{a_{\chi^2}}$, which is (III.20.3).

Rank five. A nonzero full minor including the all-ones row has absolute value at most three. The row-lattice index is therefore one, two, or three. Index one gives zero. At index two the unique nonzero relation modulo two has even support; distinctness excludes support two, so its support is four. The remaining vertex has a forced trivial character and the moment vanishes.

At index three, take a nonzero relation modulo three. Its coefficients are $0, 1, -1$. It cannot have support at most three, by the primitive three-vertex property. If it had support four, the all-ones row would force two plus signs and two minus signs. In each binary coordinate the corresponding integer sum lies in $[-2, 2]$; divisibility by three would force the sum to be zero as an integer. This would be a real parallelogram relation, contradicting rank five. Hence the relation has full support.

The all-ones row now forces four coefficients of one sign and one of the other. After rescaling the relation modulo three, write them as $(-1, 1, 1, 1, 1)$, and move the exceptional vertex to zero.

Each coordinate of the other four vertices has either zero or three ones. Thus its pattern is 0000 or the complement of a unit vector. Rank five requires all four nonzero pattern types. Grouping equal patterns gives (III.20.2), with four nonempty disjoint blocks. The constraint tuples are $(\chi^{-1}, \chi, \chi, \chi, \chi)$ with $\chi^3 = 1$. Their Fourier products give (III.20.4). $\square$

Proposition III.20.2 (Counting the five-vertex types). *The numbers of star and ternary-torsion subsets are respectively*

$$Q_d = \frac{2^d}{6}(4^d - 3 \cdot 3^d + 3 \cdot 2^d - 1), \quad (\text{III.20.5})$$

$$S_d = \frac{2^d}{24}(5^d - 4 \cdot 4^d + 6 \cdot 3^d - 4 \cdot 2^d + 1). \quad (\text{III.20.6})$$

Proof. For a star, the coefficient- -2 vertex is intrinsically distinguished. Choose this vertex and then choose three nonempty active coordinate blocks. Only their $3!$ orders produce multiple counting. This gives (III.20.5).

For a ternary configuration the exceptional coefficient in its unique modulo-three relation is distinguished. Choose that vertex and four nonempty active blocks, dividing by $4!$. Inclusion-exclusion gives (III.20.6). Unlike the binary-torsion case, one does not divide by the number of vertices: the exceptional vertex is unique. $\square$

The Stirling expressions (F.14.5) follow without any algebraic guessing. Assign d labeled coordinates to k nonempty active blocks and one possibly empty inactive block. Adjoin one extra labeled element to the inactive block. After forgetting the names of the active blocks, this is a partition of $d + 1$ elements into $k + 1$ nonempty blocks, with the block of the extra element distinguished. The number of assignments is therefore $k! \left\{ \begin{matrix} d+1 \\ k+1 \end{matrix} \right\}$.

[write] **6 October 2026; sources 06–10.** The star moment (III.20.3) is source 04's $B_5(f)$ ((III.3.8)) and source 08's $S_5(f)$; the ternary moment (III.20.4) is source 03's (III.2.29). This classification answers the fifth-order part of source 03's Research question on the fifth and sixth coefficients (note there).

III.21 Exact expansion and quantitative consequences

Source 09, delivered Section 7.

Define

$$E_4(f) = \sum_{\chi} |a_{\chi}|^4, \quad (\text{III.21.1})$$

$$T_2(f) = \sum_{\chi^2=1} |a_{\chi}|^4, \quad (\text{III.21.2})$$

$$M_5(f) = \sum_{\chi} |a_{\chi}|^2 a_{\chi}^2 \overline{a_{\chi^2}}, \quad (\text{III.21.3})$$

$$T_{3,5}(f) = \sum_{\chi^3=1} a_{\chi}^4 \overline{a_{\chi}}. \quad (\text{III.21.4})$$

The square in a_{χ}^2 means the square of a Fourier coefficient; the subscript in a_{χ^2} means evaluation at the squared character. These operations must not be confused. Both quintic statistics are real: conjugation permutes the summands by $\chi \mapsto \chi^{-1}$. They need not be nonnegative.

Theorem III.21.1 (Exact fourth and fifth orders). *Let G be finite abelian, $d \geq 3$, $m = 2^d$, $\delta \geq 0$, and $f : G \rightarrow \mathbb{R}$ have mean zero. With $\varepsilon = \|f\|_{U^d}$,*

$$\begin{aligned} \mathcal{C}_d(\delta + f) &= \delta^m + \delta^{m-4}(P_d \mathbf{E}_4(f) + T_d \mathbf{T}_2(f)) \\ &\quad + \delta^{m-5}(Q_d \mathbf{M}_5(f) + S_d \mathbf{T}_{3,5}(f)) + R_6, \end{aligned} \quad (\text{III.21.5})$$

where

$$R_6 = \sum_{\substack{S \subseteq V_d \\ |S| \geq 6}} \delta^{m-|S|} T_S(f), \quad |R_6| \leq B_{m,6}(\delta, \varepsilon).$$

The coefficients are given by (F.14.5) and satisfy $P_d + T_d = A_d$ and $Q_d = 4T_d$.

Proof. Start with (III.17.3). Apply Lemma III.17.1, Theorems III.19.1 and III.20.1, and the counting Propositions III.19.2 and III.20.2. The remaining subsets are exactly those defining R_6 . Apply Lemma III.16.1 to each one. The two coefficient identities follow from (III.19.1)–(III.20.6). $\square$

[write] **6 October 2026; sources 15–20.** Two further independent derivations: source 16’s Theorem III.39.1 (quintic part stated for odd order) and source 18’s Theorem III.40.1 (every finite abelian group, with a Smith-form proof that only 2 and 3 enter through order five). Neither source could see this one.

d	P_d	T_d	Q_d	S_d
2	1	0	0	0
3	12	2	8	0
4	100	40	160	16
5	720	520	2080	480
6	4816	5600	22400	8960
7	30912	54432	217728	134400

III.21.1 Useful special cases

For three-dimensional cubes,

$$\mathcal{C}_3(\delta + f) = \delta^8 + \delta^4(12\mathbf{E}_4 + 2\mathbf{T}_2) + 8\delta^3\mathbf{M}_5 + R_6, \quad (\text{III.21.6})$$

with

$$|R_6| \leq 28\delta^2\varepsilon^6 + 8\delta\varepsilon^7 + \varepsilon^8.$$

For four-dimensional cubes the fourth and fifth corrections are

$$\delta^{12}(100\mathbf{E}_4 + 40\mathbf{T}_2), \quad \delta^{11}(160\mathbf{M}_5 + 16\mathbf{T}_{3,5}).$$

If $|G|$ is odd, then $\mathbf{T}_2 = 0$. If $3 \nmid |G|$, then $\mathbf{T}_{3,5} = 0$. These vanishings include the trivial character because $a_1 = 0$. If G has exponent two, every $\chi^2 = 1$, so $\mathbf{M}_5 = 0$ and $\mathbf{T}_2 = \mathbf{E}_4$; there is also no ternary term. If G has exponent three, then $\chi^2 = \chi^{-1}$ and $\mathbf{M}_5 = \mathbf{T}_{3,5}$.

III.21.2 A purely Gowers-norm upper bound

Corollary III.21.2. *Under the hypotheses of Theorem III.21.1,*

$$\mathcal{C}_d(\delta + f) \leq \delta^m + A_d \delta^{m-4} \varepsilon^4 + (Q_d + S_d) \delta^{m-5} \varepsilon^5 + B_{m,6}(\delta, \varepsilon). \quad (\text{III.21.7})$$

If $|G|$ is coprime to six, A_d can be replaced by P_d and $Q_d + S_d$ by Q_d .

Proof. Each nonzero type is a subset moment and is bounded in absolute value by ε to the size of its support. Sum over the counted subsets. For a coprime-to-six group the torsion terms vanish. $\square$

For an indicator with $\varepsilon^m \leq \alpha$, this gives the explicit replacement for (F.14.1) obtained by multiplying

$$\delta^m + A_d \delta^{m-4} \alpha^{4/m} + (Q_d + S_d) \delta^{m-5} \alpha^{5/m} + \sum_{s=6}^m \binom{m}{s} \delta^{m-s} \alpha^{s/m}$$

by N^{d+1} . The separate exact formula (III.17.4) applies when $d = 2$.

III.21.3 Keeping Fourier information instead of discarding it

Let

$$\tau = \|f\|_{U^2}, \quad \rho = \max_{\chi} |a_{\chi}|, \quad \mathsf{T}_{3,4}(f) = \sum_{\chi^3=1} |a_{\chi}|^4.$$

Then $\rho \leq \tau \leq \varepsilon$, and directly from the definitions,

$$|\mathsf{M}_5| \leq \rho \tau^4, \quad |\mathsf{T}_{3,5}| \leq \rho \mathsf{T}_{3,4} \leq \rho \tau^4. \quad (\text{III.21.8})$$

Consequently,

$$\begin{aligned} \mathcal{C}_d(\delta + f) - \delta^m &\leq \delta^{m-4} (P_d \tau^4 + T_d \mathsf{T}_2) \\ &\quad + \delta^{m-5} \rho (Q_d \tau^4 + S_d \mathsf{T}_{3,4}) + B_{m,6}(\delta, \varepsilon). \end{aligned} \quad (\text{III.21.9})$$

This is useful when ordinary Fourier uniformity is substantially stronger than the available higher-order uniformity. The exact formula shows that all terms through degree five are ordinary Fourier statistics; one should not pay a higher-order uniformity bound for them unnecessarily.

III.21.4 Certificates of nonuniformity

For $F : G \rightarrow [0, 1]$, put $\Delta = \mathcal{C}_d(F) - \delta^m \geq 0$. Since $|f| \leq 1$, one has $\varepsilon \leq 1$. Therefore

$$\varepsilon \geq \left[\frac{\Delta}{\binom{m}{4} (1 + \delta)^{m-4}} \right]^{1/4}. \quad (\text{III.21.10})$$

This follows from Theorem III.17.2 and (III.16.5). A sharper, still explicit denominator follows from (III.21.7) by replacing every ε^s , $s \geq 4$, by ε^4 .

If the fourth and fifth Fourier corrections have been measured, define

$$D_6 = \mathcal{C}_d(F) - \delta^m - \delta^{m-4} (P_d \mathsf{E}_4 + T_d \mathsf{T}_2) - \delta^{m-5} (Q_d \mathsf{M}_5 + S_d \mathsf{T}_{3,5}).$$

Then

$$\varepsilon \geq \left[\frac{|D_6|}{\binom{m}{6}(1+\delta)^{m-6}} \right]^{1/6} \quad (d \geq 3). \quad (\text{III.21.11})$$

These are lower bounds on a uniformity norm. They do not, by themselves, produce a polynomial phase, a nilsequence, or a density increment. Calling them inverse theorems would overstate their content.

III.21.5 Local derivatives at a constant

For a fixed mean-zero direction f , the map $t \mapsto \mathcal{C}_d(\delta + tf)$ is a polynomial. Its first three derivatives at zero vanish, and

$$\begin{aligned} \left. \frac{d^4}{dt^4} \mathcal{C}_d(\delta + tf) \right|_{t=0} &= 4! \delta^{m-4} (P_d \mathbf{E}_4(f) + T_d \mathbf{T}_2(f)), \\ \left. \frac{d^5}{dt^5} \mathcal{C}_d(\delta + tf) \right|_{t=0} &= 5! \delta^{m-5} (Q_d \mathbf{M}_5(f) + S_d \mathbf{T}_{3,5}(f)). \end{aligned}$$

Thus the cube-count functional is flat through third order along the mean-zero hyperplane at a positive constant. Its first nonzero variation is quartic, and its arithmetic anisotropy is already visible there. This is a statement about a specific functional at a specific base point, not a claim of vanishing Hessians at arbitrary functions.

III.22 Sharpness, signs, and transfer to genuine sets

Source 09, delivered Section 8.

III.22.1 The fourth-order exponent and its universal coefficient

Proposition III.22.1. *For every $d \geq 2$ and $0 < \delta < 1$, no uniform replacement of the fourth-order error by $O_{d,\delta}(\varepsilon^{4+\eta})$, $\eta > 0$, is possible. This remains true when groups are restricted to have odd order. If all finite abelian groups are allowed, the coefficient A_d of $\delta^{m-4}\varepsilon^4$ in (III.21.7) is optimal as a leading small- ε coefficient.*

Proof. Fix any nonzero real mean-zero function g on a finite group and put $F_t = \delta + tg$, with $t > 0$ small enough that $0 \leq F_t \leq 1$. Its balanced norm is $t\|g\|_{U^d}$. The fourth-order coefficient in Theorem III.21.1 is positive because $P_d > 0$ and $\mathbf{E}_4(g) > 0$. For $d = 2$, use (III.17.4) instead. Thus the excess is a positive constant times $t^4 + O(t^5)$ and cannot be of order $t^{4+\eta}$. A nonzero cosine on a cyclic group of odd prime order at least seventeen gives the same conclusion among odd-order groups.

For the coefficient assertion, take a nontrivial order-two character $\chi : G \rightarrow \{\pm 1\}$ and let $g = \chi$. Then $\|g\|_{U^d} = 1$, $\mathbf{E}_4(g) = \mathbf{T}_2(g) = 1$, and the leading coefficient is exactly $P_d + T_d = A_d$. This rules out a smaller universal leading coefficient. $\square$

There is also a closed formula for the last example:

$$\mathcal{C}_d(\delta + t\chi) = 2^{-d-1}((\delta + t)^m + (\delta - t)^m) + (1 - 2^{-d})(\delta^2 - t^2)^{m/2}. \quad (\text{III.22.1})$$

To prove it, condition on $(\chi(h_1), \dots, \chi(h_d))$. With probability 2^{-d} all signs are positive and every vertex has sign $\chi(x)$, whose two values are equiprobable. Otherwise exactly half the cube vertices have each sign. In dimension three, (III.22.1) reduces to

$$\mathcal{C}_3(\delta + t\chi) = \delta^8 + 14\delta^4 t^4 + t^8.$$

This also explains why deleting the binary-torsion term would give an incorrect fourth-order coefficient.

III.22.2 A negative quintic correction

Proposition III.22.2. *The fifth-order term in Theorem III.21.1 can have either sign, even on groups of prime order greater than three. In particular, the inequality*

$$\mathcal{C}_3(\delta + f) \geq \delta^8 + 12\delta^4 \|f\|_{U^2}^4 \quad (\text{III.22.2})$$

is false in general, even for $[0, 1]$ -valued functions on such groups.

Proof. Let $q \geq 17$ be prime and, for $\sigma \in \{1, -1\}$, put

$$g_\sigma(x) = \cos(2\pi x/q) + \sigma \cos(4\pi x/q).$$

Its only nonzero Fourier coefficients are $a_{\pm 1} = 1/2$ and $a_{\pm 2} = \sigma/2$, with frequencies written additively in this calculation. Therefore

$$\mathsf{E}_4(g_\sigma) = \frac{1}{4}, \quad \mathsf{M}_5(g_\sigma) = \frac{\sigma}{16}, \quad \mathsf{T}_2(g_\sigma) = \mathsf{T}_{3,5}(g_\sigma) = 0.$$

Apply (III.21.6) to $f = tg_\sigma$. At any fixed $0 < \delta < 1$,

$$\mathcal{C}_3(\delta + tg_\sigma) = \delta^8 + 3\delta^4 t^4 + \frac{\sigma}{2} \delta^3 t^5 + O_\delta(t^6). \quad (\text{III.22.3})$$

For all sufficiently small $t > 0$, one has $0 \leq \delta + tg_\sigma \leq 1$ because $|g_\sigma| \leq 2$. Taking $\sigma = -1$ makes the difference between the left side of (III.22.2) and its proposed right side negative. The $O(t^6)$ bound is explicitly supplied by (F.14.6); no finite numerical experiment is needed for this counterexample. $\square$

The same example proves that, after the exact fourth-order correction has been removed, a fifth-order remainder cannot in general be replaced by a higher power. More strongly, for every $d \geq 3$ the star count Q_d is positive, so the same two-frequency example gives a nonzero fifth-order term in that dimension.

An exact finite calculation, reproduced in Appendix III.26, yields

$$\mathcal{C}_3(\delta + tg_\sigma) = \delta^8 + 3\delta^4 t^4 + \frac{\sigma}{2} \delta^3 t^5 + \frac{3}{2} \delta^2 t^6 + \frac{1}{4} t^8. \quad (\text{III.22.4})$$

For example, when $0 < t \leq \delta/4$, the negative fifth-order term in the $\sigma = -1$ case dominates the displayed positive sixth- and eighth-order terms. The asymptotic proof above is logically independent of this stronger computational identity.

[write] **6 October 2026; sources 06–10.** (III.22.2) is the bound refuted also by source 03 (on $\mathbb{Z}/3\mathbb{Z}$) and source 08 (on $\mathbb{Z}/11\mathbb{Z}$, Theorem III.10.1); Remark W1 keeps it on record with an explicit counterexample. The proposition’s “groups of prime order greater than three” is witnessed by the proof for primes $q \geq 17$; on $\mathbb{Z}/11\mathbb{Z}$ source 08’s example is the case $\sigma = -1$ and its polynomial is (III.22.4). On $\mathbb{Z}/5\mathbb{Z}$ the fifth-order argument does not apply to g_{-1} : there $2 \cdot 2 = -1$, the coefficients at ± 2 also contribute, and $\mathsf{M}_5(g_{-1}) = 0$ (checked at the second write).

III.22.3 The sixth-order remainder is also necessary

Proposition III.22.3. *For every $d \geq 3$ and $0 < \delta < 1$, the sixth-order remainder after the exact fourth and fifth corrections cannot be uniformly replaced by $O_{d,\delta}(\varepsilon^{6+\eta})$ for any $\eta > 0$.*

Proof. Use g_+ from Proposition III.22.2. Every Fourier coefficient is nonnegative, so every subset moment is nonnegative by Proposition III.18.1. Consider the six vertices of a three-dimensional cube obtained by deleting the adjacent pair $(0, 0, 0), (1, 0, 0)$. On the three remaining vertices in the first slice assign frequencies $(1, 1, -2)$, and on the corresponding vertices in the second slice assign $(-1, -1, 2)$. Their total frequency is zero; the sum in the first cube coordinate is zero; and the sums in each of the other two coordinates cancel between the slices. This is an admissible Fourier tuple with positive coefficient product. Thus that six-vertex moment is strictly positive.

For $d > 3$, embed this subset in a three-dimensional face and give the remaining coordinates value zero. The same tuple still satisfies all constraints. Hence the coefficient of t^6 in R_6 for $f = tg_+$ is positive, and no other sixth-order contribution can cancel it. Since $\varepsilon = t\|g_+\|_{U^d}$, the claimed stronger remainder is impossible. $\square$

III.22.4 Why weighted examples are not a loophole

The following rounding lemma transfers these phenomena to indicators. It also quantifies the size of the auxiliary group needed for the transfer.

Lemma III.22.4 (Random rounding in a Gowers norm). *Let $F : G \rightarrow [0, 1]$, let $d \geq 2$, and let $m = 2^d$. There exists $A \subseteq G$ such that*

$$\|\mathbf{1}_A - F\|_{U^d} \leq \zeta := \left(\frac{\binom{m}{2}}{|G|} \right)^{1/m}. \quad (\text{III.22.5})$$

Writing $\delta = \mathbb{E}F$ and $\delta_A = |A|/|G|$, one also has

$$|\delta_A - \delta| \leq \zeta, \quad (\text{III.22.6})$$

$$|\mathcal{C}_d(\mathbf{1}_A) - \mathcal{C}_d(F)| \leq m\zeta, \quad (\text{III.22.7})$$

$$\|(\mathbf{1}_A - \delta_A) - (F - \delta)\|_{U^d} \leq 2\zeta. \quad (\text{III.22.8})$$

Proof. Choose independent Bernoulli variables $\mathbf{1}_A(x)$ of means $F(x)$, and set $R(x) = \mathbf{1}_A(x) - F(x)$. For a fixed parameterized cube whose m evaluated vertices are distinct, independence and $\mathbb{E}_A R(x) = 0$ give $\mathbb{E}_A \prod_v R(L_v) = 0$. If vertices collide, the absolute value of the expected product is at most one.

For two distinct abstract vertices v, w , the equation $L_v = L_w$ has probability $1/|G|$: one coefficient in $\sum_i (v_i - w_i)h_i$ is 1 or -1 . The union bound gives collision probability at most $\binom{m}{2}/|G|$. Therefore

$$0 \leq \mathbb{E}_A \|R\|_{U^d}^m \leq \frac{\binom{m}{2}}{|G|}.$$

Some realization satisfies (III.22.5). Now $|\mathbb{E}R| \leq \|R\|_{U^d}$ proves (III.22.6). Telescope the two cube products in (III.22.7); each of the m terms has one factor R , and all other factors are either F or $\mathbf{1}_A$, with Gowers norm at most one. Lemma III.16.1 bounds each term by ζ . Finally, subtract the mean of R and apply the triangle inequality to prove (III.22.8). $\square$

Corollary III.22.5 (Transfer of the sharpness examples). *The fourth-order exponent obstruction, the optimal universal leading coefficient A_d , the fifth-order sign obstruction, and the necessity of a sixth-order remainder persist for sequences of actual sets in growing finite abelian groups. Their densities tend to any prescribed $\delta \in (0, 1)$. The first exponent obstruction also persists in odd-order groups.*

Proof. Lift a fixed weighted example on H to $H \times K$ using the projection to H . All cube averages and Gowers norms of the lifted function are unchanged. Let $t \rightarrow 0$ and choose $|K|$ sufficiently large that the rounding error satisfies $\zeta = o(t^6)$; larger polynomial choices of $|K|$ suffice.

The balanced norm changes by at most 2ζ , the density by at most ζ , and the full cube count by at most $m\zeta$. Each of the fourth- and fifth-order statistics is itself a fixed subset moment. Telescoping it and using Gowers–Cauchy–Schwarz shows that it changes by $O_d(\zeta)$, since the balanced functions involved are bounded by one. Thus the nonzero terms at orders four, five, and six persist. For the odd-order assertion take both H and K of odd order. $\square$

Remark III.22.6 (Exact rational densities). For a prescribed rational $\delta = a/b \in (0, 1)$, the same obstructions can be realized with density exactly δ . Choose $|H \times K|$ divisible by b . After rounding, add or delete at most $\zeta|G|$ points to obtain the required cardinality. If E is the edit set, the difference D of the two indicators is supported on E and bounded by one. Hence

$$\|D\|_{U^d}^m \leq \mathbb{E}_{x,h} \left| \prod_v D(L_v) \right| \leq \frac{|E|}{|G|} \leq \zeta.$$

The total approximation error in U^d is at most $\zeta + \zeta^{1/m}$. Choosing $\zeta = o(t^{6m})$ makes this $o(t^6)$ and preserves all the sharpness conclusions at the exact fixed density. For an odd-order construction this exact-density variant applies when b is odd; arbitrary densities can still be approached by the preceding corollary. Irrational densities, of course, cannot be the exact density of a finite set.

[write] **6 October 2026; sources 06–10.** Rounding to sets is also done by sources 03, 04 and 08 (note after source 08’s Lemma III.11.2). Lemma III.22.4 has constant $((\binom{m}{2})/|G|)^{1/m}$ by taking a realization at most the mean; source 08 takes twice the mean.

III.23 A sixth-order order-five obstruction and an all-prime family

Source 09, delivered Section 9.

The small-support expansion is especially simple on groups whose order is coprime to six. That simplification does not continue indefinitely. There is already an order-five torsion configuration on six selected vertices.

Theorem III.23.1 (First possible isolated order-five contribution). *There exists a six-vertex subset of V_5 with full augmented rank and row-lattice index five, whose mean-zero moment can be nonzero on groups with order-five characters but is zero when multiplication by five is invertible on G . No such order-five-only contribution occurs on fewer than six selected cube vertices.*

Proof. Consider the binary matrix

$$B = \begin{pmatrix} 1 & 0 & 1 & 1 & 1 \\ 0 & 1 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 0 & 1 \\ 1 & 1 & 0 & 1 & 0 \end{pmatrix}. \quad (\text{III.23.1})$$

Its determinant is -5 . Take S to consist of the zero vector and the five columns of B . Its augmented matrix is

$$M_S = \begin{pmatrix} 1 & \mathbf{1}^T \\ 0 & B \end{pmatrix}, \quad |\det M_S| = 5.$$

For

$$r = (2, 1, 3, 2, 1, 1)^T$$

one checks exactly that

$$M_S r = (10, 5, 5, 5, 5, 5)^T. \quad (\text{III.23.2})$$

Every coordinate of r is nonzero modulo five. Since the row quotient has order five, the full character constraint group is

$$K_G(S) = \{(\chi^{r_0}, \dots, \chi^{r_5}) : \chi^5 = 1\}.$$

Consequently,

$$T_S(f) = \sum_{\chi^5=1} \prod_{j=0}^5 a_{\chi^{r_j}}. \quad (\text{III.23.3})$$

If multiplication by five is invertible on G , the only parameter is the trivial character, and the moment is zero because $a_1 = 0$. On $G = \mathbb{Z}/5\mathbb{Z}$, take $f = 5\mathbf{1}_{\{0\}} - 1$. Its trivial Fourier coefficient is zero and its four nontrivial Fourier coefficients are one. Thus (III.23.3) is exactly four.

Finally, Lemma III.17.1 and the complete four- and five-vertex classifications show that the only isolated torsion contributions before degree six have orders two and three. In particular, none can be supported purely on order-five characters. This proves the minimal support-size assertion. $\square$

The determinant in (III.23.1) can be checked by ordinary integer elimination; the accompanying exact script additionally records the Smith invariants $1, 1, 1, 1, 1, 5$ of M_S . A direct elimination certificate is given in Appendix III.27. Rescaling the displayed f by a sufficiently small t and adding a density gives a $[0, 1]$ -valued example, so this obstruction belongs to the same perturbative counting problem as the main theorem.

[write] **6 October 2026; sources 06–10.** *Checked at the second write* with SymPy 1.14.0: $\det B = -5$, $\det M_S = -5$, $M_S r = (10, 5, 5, 5, 5, 5)$, the Smith normal form of M_S is $\text{diag}(1, 1, 1, 1, 1, 5)$, the five columns of B are distinct nonzero binary vectors, and the moment (III.23.3) equals 4 for $f = 5\mathbf{1}_{\{0\}} - 1$ on $\mathbb{Z}/5\mathbb{Z}$. At four vertices only the invariant factors one and two occur (source 03's Remark III.2.4), at five also three (Theorem III.20.1); index five first occurs here, at six vertices.

Proposition III.23.2 (An isolated torsion configuration for every prime). *For every prime ℓ , there is a subset of $V_{\ell+1}$ with $\ell + 2$ vertices, full augmented rank, and row-lattice quotient $\mathbb{Z}/\ell\mathbb{Z}$, whose constraint relation has full support modulo ℓ .*

Proof. Put $n = \ell + 1$, and take the zero vector together with the n columns of $J - I_n$. The eigenvalues of $J - I_n$ are $n - 1 = \ell$ on the all-ones direction and -1 on its complementary hyperplane. Hence its determinant is $(-1)^{n-1}\ell$. Its augmented matrix has the same determinant and therefore a quotient of order ℓ . The relation modulo ℓ is

$$(-1, 1, \dots, 1),$$

because its sum is $n - 1 = \ell$ and each coordinate row has ℓ ones. Every coefficient is nonzero. The moment is

$$\sum_{\chi^\ell=1} a_\chi^{\ell+1} \overline{a_\chi}.$$

It is nonzero for $f = \ell \mathbf{1}_{\{0\}} - 1$ on $\mathbb{Z}/\ell\mathbb{Z}$. □

For $\ell = 2, 3$ this reproduces the two small torsion normal forms. For $\ell = 5$ it uses seven vertices, whereas Theorem III.23.1 uses six. Thus the uniform construction is not generally support-minimal.

One can make the resulting extremal problem precise. Let $h(\ell)$ be the least cardinality of an affinely independent Boolean-vertex set whose integer row quotient has ℓ -torsion and whose character kernel over $\mathbb{Z}/\ell\mathbb{Z}$ has a vector nonzero in every coordinate. The results above give

$$h(2) = 4, \quad h(3) = 5, \quad h(5) = 6, \quad h(\ell) \leq \ell + 2.$$

If $s = h(\ell)$, some nonzero full minor including the all-ones row is divisible by ℓ . The sign-matrix argument in Lemma III.18.2 gives the necessary condition

$$\ell \leq \frac{s^{s/2}}{2^{s-1}}. \tag{III.23.4}$$

In particular, as ℓ grows, $h(\ell) \geq (2 - o(1)) \log \ell / \log \log \ell$. There is a wide gap between this elementary obstruction and the explicit general construction. We leave its systematic improvement as a research problem rather than pretending the small examples solve it.

III.24 Heterogeneous functions and an all-orders organizing principle

Source 09, delivered Section 10.

The method is not restricted to a single set placed at every vertex. Let $F_v = \delta_v + f_v$ be real functions indexed by V_d , with $\mathbb{E}f_v = 0$, $\delta_v \geq 0$, and $\varepsilon_v = \|f_v\|_{U^d}$.

Theorem III.24.1 (Heterogeneous expansion). *One has*

$$\begin{aligned} \mathbb{E}_{x,h} \prod_{v \in V_d} F_v(L_v) &= \prod_{v \in V_d} \delta_v \\ &+ \sum_{\substack{S \subseteq V_d \\ |S| \geq 4}} \left(\prod_{v \notin S} \delta_v \right) \mathbb{E}_{x,h} \prod_{v \in S} f_v(L_v). \end{aligned} \tag{III.24.1}$$

Every four- or five-vertex term vanishes unless its vertex subset is one of the normal forms in Theorems III.19.1 and III.20.1. For a surviving term its exact value is the mixed Fourier sum in

(III.18.3). *The remainder after size five is bounded by*

$$\sum_{\substack{S \subseteq V_d \\ |S| \geq 6}} \left(\prod_{v \notin S} \delta_v \right) \left(\prod_{v \in S} \varepsilon_v \right). \quad (\text{III.24.2})$$

Proof. Expand the product. Joint uniformity for up to three selected vertices makes those mixed moments a product of zero means. In the classification proofs, every vanishing case either has fully uniform values or forces a trivial character at one selected vertex, so the same vanishing holds for different mean-zero functions. Finally, apply Lemma III.16.1 to bound each remaining mixed term. $\square$

For example, a parallelogram with ordered relation $v_1 + v_4 = v_2 + v_3$ contributes

$$\sum_{\chi} a_{1,\chi} a_{4,\chi} a_{2,\chi^{-1}} a_{3,\chi^{-1}}.$$

A binary-torsion subset contributes $\sum_{\chi^2=1} \prod_{j=1}^4 a_{j,\chi}$. Thus the homogeneous formulas are not isolated coincidences; they are the diagonal specializations of a colored cube-count expansion.

At every order, Proposition III.18.1 gives a finite algorithm: enumerate a support S , compute its integer row lattice, identify which character constraints admit full-support tuples, and sum the corresponding mixed Fourier products. The difficult part is not the existence of an algorithm but the compression of the support types into dimension-uniform, human-readable formulas. Through degree five, the four normal forms achieve precisely that compression.

III.25 Relation to progression bounds and current conjecture status

Source 09, delivered Section 13.

The local result should be placed against the right global benchmark. Gowers' 2001 theorem is not the current quantitative endpoint. For example, Leng, Sah, and Sawhney obtain, for every fixed $k \geq 5$, an upper bound of the form

$$r_k(N) \ll N \exp(-(\log \log N)^{c_k})$$

with $c_k > 0$ [2]. The present manuscript does not compete with that bound: it sharpens one local cube-count inequality, without supplying a new inverse theorem or a new iteration.

A recent primary preprint of Milićević establishes a quasipolynomial inverse theorem for $U^k(\mathbb{F}_p^n)$ in the high characteristic regime $p \geq k$ [32]. That is another reason not to describe an elementary Fourier expansion as a breakthrough inverse theorem. Its relation to this manuscript is potentially complementary: low-order terms can be evaluated before an inverse theorem is applied to a residual, but a useful quantitative implementation would need a separate argument.

The conjectural lower bounds in Section 4 of Gowers' paper also require an up-to-date reading. Shi and Dong's 2026 preprint states a disproof of Gowers' conjectured progression lower bound for every even length at least six, and records the earlier length-four disproof [33]. Their higher-length uniformity condition is U^{k-2} uniformity, not merely ordinary Fourier uniformity. We cite this as the

stated result of that preprint, rather than treating the old conjecture as still untouched or claiming its resolution here.

The cube lower bound $\mathcal{C}_d(F) \geq \delta^m$ is different. It follows from a norm inequality and remains valid. Proposition III.22.2 does not refute that bound: it refutes the stronger attempt to add the entire positive quartic correction as a global lower bound while ignoring the signed quintic term. Arithmetic progression counts and full cube counts are different linear-configuration averages; one must not transfer extremal conclusions between them without proof.

III.26 Exact two-frequency polynomial certificate

Source 09, delivered Appendix A.

For $q \geq 17$, all character constraints for eight frequencies from $\{-2, -1, 0, 1, 2\}$ can be checked as integer equalities: the total absolute frequency sum is at most sixteen, and every coordinate sum has absolute value at most eight. Thus no nonzero multiple of q can occur.

Order the eight vertices $v \in \{0, 1\}^3$. A term in the expansion of $\mathcal{C}_3(\delta + tg_\sigma)$ is indexed by frequencies $b_v \in \{-2, -1, 0, 1, 2\}$ satisfying

$$\sum_v b_v = 0, \quad \sum_v v_i b_v = 0 \quad (i = 1, 2, 3). \quad (\text{III.26.1})$$

Let s be the number of nonzero frequencies and r the number with absolute value two. Its contribution is $\delta^{8-s} t^s 2^{-s} \sigma^r$. The exact signed counts are

	s	0	1	2	3	4	5	6	7	8
$\sum 1$	1	1	0	0	0	48	16	96	0	64
$\sum (-1)^r$	1	1	0	0	0	48	-16	96	0	64

Multiplying the s th column by $2^{-s} \delta^{8-s} t^s$ proves (III.22.4). The following finite loop specifies the certificate without a floating-point operation:

```
for b in product(range(-2, 3), repeat=8):
    if sum(b) != 0:
        continue
    if any(sum(b[j] for j in coordinate_slice[i]) != 0
           for i in range(3)):
        continue
    s = sum(x != 0 for x in b)
    r = sum(abs(x) == 2 for x in b)
    positive[s] += 1
    signed[s] += (-1)**r
```

The complete executable version is included in `verify_expansion.py`, along with assertions for every table entry. This is a transparent exhaustive finite certificate; the all-dimension expansion and its sharpness proofs do not depend on trusting a numerical approximation to this table.

III.27 An integer determinant certificate for the order-five example

Source 09, delivered Appendix B.

Here is a direct check of the determinant in (III.23.1). Subtract row one from rows four and five, and expand along column one. The remaining 4×4 matrix is

$$C = \begin{pmatrix} 1 & 1 & 0 & 0 \\ 1 & 0 & 1 & 1 \\ 1 & -1 & -1 & 0 \\ 1 & -1 & 0 & -1 \end{pmatrix}.$$

Subtract its first row from each remaining row and expand again along column one. This gives

$$D = \begin{pmatrix} -1 & 1 & 1 \\ -2 & -1 & 0 \\ -2 & 0 & -1 \end{pmatrix}, \quad \det D = -1 - 2 - 2 = -5.$$

All these row operations preserve the determinant, and both expansions use the upper-left entry with positive cofactor sign. Thus $\det B = -5$. Together with (III.23.2), this proves the index and kernel claims without reliance on the Smith-normal-form implementation.

III.28 A torsion-sensitive generalized von Neumann inequality

Source 10, delivered Section 5.

The reindexing step in the proof of Theorem 3.2 of the source paper is where the ambient modulus enters. Retaining the sizes of the relevant fibres gives an explicit estimate on every finite abelian group. Retaining the norms of the other functions at the same time gives a useful density improvement in the subsequent counting argument.

Throughout this section G is a nonempty finite abelian group, $N = |G|$, and all averages and L^p norms are normalized. For an integer c put

$$G[c] = \{x \in G : cx = 0\}, \quad \kappa_G(c) = |G[c]|.$$

Thus $\kappa_G(0) = N$, and multiplication by c is invertible exactly when $\kappa_G(c) = 1$. On a cyclic group, $\kappa_{\mathbb{Z}/N\mathbb{Z}}(c) = \gcd(N, c)$, with the usual convention $\gcd(N, 0) = N$.

For $f : G \rightarrow \mathbb{C}$ set $\Delta_h f(x) = f(x) \overline{f(x+h)}$, and write

$$\|f\|_{U^s}^{2^s} = \mathbb{E}_{x, h_1, \dots, h_s \in G} \prod_{\omega \in \{0,1\}^s} \mathcal{C}^{|\omega|} f(x + \omega \cdot h), \quad s \geq 1,$$

where $\mathcal{C}$ denotes complex conjugation. These powers are nonnegative. In particular $\|f\|_{U^1} = |\mathbb{E}f|$, which is a seminorm, and for $s \geq 2$ the recursion is

$$\|f\|_{U^s}^{2^s} = \mathbb{E}_h \|\Delta_h f\|_{U^{s-1}}^{2^{s-1}}. \quad (\text{III.28.1})$$

Lemma III.28.1 (The exact fibre identity). *For every $V : G \rightarrow [0, \infty)$ and every integer c ,*

$$\mathbb{E}_u V(cu) = \kappa_G(c) \mathbb{E}_h \mathbf{1}_{cG}(h) V(h) \leq \kappa_G(c) \mathbb{E}_h V(h).$$

Proof. Every element of the image cG has exactly $\kappa_G(c)$ preimages. Counting those preimages proves the equality, and $V \geq 0$ proves the inequality. $\square$

For a coefficient vector $\mathbf{a} = (a_1, \dots, a_k)$ define

$$\Lambda_{\mathbf{a}}(f_1, \dots, f_k) = \mathbb{E}_{x,d} \prod_{i=1}^k f_i(x + a_i d).$$

Theorem III.28.2 (General coefficients and mixed L^p control). *Let $k \geq 2$, let $a_1, \dots, a_k$ be integers, and suppose that multiplication by $a_k - a_{k-1}$ is invertible on G . Set*

$$K = \prod_{i=1}^{k-1} \kappa_G(a_k - a_i), \quad M = \left(\prod_{i=1}^{k-2} \|f_i\|_{L^{2^i}} \right) \|f_{k-1}\|_{L^{2^{k-2}}}.$$

For arbitrary functions $f_i : G \rightarrow \mathbb{C}$,

$$|\Lambda_{\mathbf{a}}(f_1, \dots, f_k)| \leq MK^{1/2^{k-1}} \|f_k\|_{U^{k-1}}. \quad (\text{III.28.2})$$

No boundedness hypothesis on the functions is needed.

Proof. For $k = 2$, the map

$$(x, d) \mapsto (x + a_1 d, x + a_2 d)$$

is a bijection of G^2 . Hence the average equals $(\mathbb{E}f_1)(\mathbb{E}f_2)$, and its absolute value is at most $\|f_1\|_1 \|f_2\|_{U^1}$, as required.

Suppose $k \geq 3$ and assume the result for $k-1$. Put $c_i = a_i - a_1$. Translate the first variable by $a_1 d$ and apply Cauchy–Schwarz in that variable. Expanding the square and writing the second difference variable as $d + h$ gives

$$\begin{aligned} |\Lambda_{\mathbf{a}}(f_1, \dots, f_k)|^2 &\leq \|f_1\|_2^2 \mathbb{E}_x \left| \mathbb{E}_d \prod_{i=2}^k f_i(x + c_i d) \right|^2 \\ &= \|f_1\|_2^2 \mathbb{E}_h \Lambda_{(c_2, \dots, c_k)}(\Delta_{c_2 h} f_2, \dots, \Delta_{c_k h} f_k). \end{aligned} \quad (\text{III.28.3})$$

The final average is real and nonnegative because it arose by expanding an average of squares. It is therefore bounded above by the average of the absolute values of its inner multilinear forms.

The induction applies to each such form: the last coefficient difference is still $a_k - a_{k-1}$. Its kernel product is

$$K' = \prod_{i=2}^{k-1} \kappa_G(a_k - a_i).$$

For every $p \geq 1$, Hölder's inequality and translation invariance give

$$\|\Delta_h f\|_{L^p} \leq \|f\|_{L^{2p}}^2. \quad (\text{III.28.4})$$

Consequently the product of the non-target norms in the induction is at most M_{rest}^2 , where

$$M_{\text{rest}} = \left(\prod_{i=2}^{k-2} \|f_i\|_{L^{2^i}} \right) \|f_{k-1}\|_{L^{2^{k-2}}}.$$

Let $q = 2^{k-2}$. A further application of Hölder, followed by Lemma III.28.1 and (III.28.1), gives

$$\begin{aligned}\mathbb{E}_h \|\Delta_{(a_k - a_1)h} f_k\|_{U^{k-2}} &\leq \left(\mathbb{E}_h \|\Delta_{(a_k - a_1)h} f_k\|_{U^{k-2}}^q \right)^{1/q} \\ &\leq \kappa_G(a_k - a_1)^{1/q} \|f_k\|_{U^{k-1}}^2.\end{aligned}$$

Substituting these bounds into (III.28.3) and taking square roots proves (III.28.2), since $M = \|f_1\|_2 M_{\text{rest}}$ and $(K' \kappa_G(a_k - a_1))^{1/(2q)} = K^{1/2^{k-1}}$. $\square$

[write] **6 October 2026; sources 06–10.** *Text repaired.* In the delivered source the subscript M_{rest} appears three times in this proof as “M_{\text{rest}}”, a line break and “m rest}”: apparently the two characters \rm of \rm were read as a carriage-return escape when the source was generated, and the delivered PDF prints “M_{m rest}”. Restored here as M_{rest} (byte scan at intake: no other control-character damage in sources 01–10).

Credits. For k -term progressions with consecutive slopes and invertible differences this is the generalized von Neumann bound of sources 03, 04 and 06 (Lemmas III.1.1, III.4.2, Theorem III.5.1) with L^∞ norms replaced by mixed L^p norms; the kernel factor $K^{1/2^{k-1}}$, which removes the invertibility assumption on all but one difference, is only here. It advances source 06’s Question 7 for one distinguished function, with explicit kernel constants that are shown sharp only for the families of Proposition III.28.4; the two-function form remains open.

[write] **6 October 2026; sources 15–20.** Source 20’s Theorem III.44.1 gives the general directional-to-global comparison with the optimal constant for each subgroup tuple, and its Proposition III.46.1 the sharp pair correction $\sqrt{m-1}$ (item III.9).

The L^p exponents assigned to the non-target functions are 1 for $k = 2$, then $(2, 2)$, $(2, 4, 4)$, $(2, 4, 8, 8)$, and so forth. Their reciprocals sum to one. The functions and coefficients can be reordered before applying the theorem, provided a coefficient with invertible difference from the target is left immediately before it. This permits optimization of the norm product when the functions have different densities or sizes.

Corollary III.28.3 (Every target in an arithmetic progression). *Let $k \geq 2$, let $f_0, \dots, f_{k-1} : G \rightarrow \mathbb{C}$ satisfy $|f_i| \leq 1$, and let $0 \leq j \leq k-1$. Then*

$$\left| \mathbb{E}_{x,d} \prod_{i=0}^{k-1} f_i(x + id) \right| \leq \left(\prod_{i \neq j} \kappa_G(i - j) \right)^{1/2^{k-1}} \|f_j\|_{U^{k-1}}. \quad (\text{III.28.5})$$

In particular, for an endpoint target the factor is $(\prod_{r=1}^{k-1} \kappa_G(r))^{1/2^{k-1}}$. It is one whenever the exponent of G is coprime to $(k-1)!$.

Proof. Every target has a neighboring coefficient with difference 1 or -1 . Place that neighbor last among the non-target functions and apply Theorem III.28.2. Its L^p product is at most one. The endpoint formula and the coprimality assertion follow directly. $\square$

For the internal target j , the kernel product is

$$\left(\prod_{r=1}^j \kappa_G(r) \right) \left(\prod_{r=1}^{k-1-j} \kappa_G(r) \right).$$

Thus coprimality of the exponent of G with $j!(k-1-j)!$ suffices for that individual estimate. The endpoint result in particular extends the prime-modulus statement to composite cyclic groups with

no prime factor below k , and to arbitrary finite abelian groups satisfying the same restriction on their prime divisors.

The invertible-neighbor hypothesis in the general-coefficient theorem is essential to the theorem as stated. For $k = 2$, coefficients $(0, 2)$ on $G = \mathbb{Z}/2\mathbb{Z}$, and opposite nontrivial characters, the multilinear average is one whereas the target U^1 seminorm is zero.

III.28.1 Sharpness and an explicit obstruction

Proposition III.28.4 (Sharp torsion growth). *For three-term progressions on $G = (\mathbb{Z}/2\mathbb{Z})^n$, the endpoint factor $|G|^{1/4}$ in (III.28.5) is exactly sharp. More generally, for $k = p + 1$ on $G = (\mathbb{Z}/p\mathbb{Z})^n$, with p prime, the endpoint factor $|G|^{1/2^p}$ has the optimal order in $|G|$, up to constants depending only on p .*

Proof. First let $G = (\mathbb{Z}/2\mathbb{Z})^n$, $N = 2^n$, and set

$$g(x_1, \dots, x_n) = i^{x_1 + \dots + x_n},$$

using representatives $x_i \in \{0, 1\}$. Every normalized Fourier coefficient of g has absolute value $N^{-1/2}$: in each coordinate the Fourier factors are $(1 + i)/2$ and $(1 - i)/2$. Therefore

$$\|g\|_{U^2}^4 = \sum_{\chi} |\hat{g}(\chi)|^4 = N^{-1}.$$

Take the two endpoint functions to be $\bar{g}, g$, and the middle function to be one. Since $2d = 0$, the progression average equals one. The upper bound is exactly $N^{1/4}N^{-1/4} = 1$.

For the general assertion write $s = p$ and $N = p^n$. Choose independent uniform unit-circle phases $g(z)$ for the points $z \in G$. For a fixed s -dimensional cube, the expectation of its alternating product is either zero or one, according as the exponents of these independent phases fail to cancel or all cancel. If its 2^s vertices are all distinct, the expectation is zero. Each pair of distinct cube vertices coincides with probability $1/N$, because their difference contains a coefficient 1 or -1 on at least one independent uniform increment. The union bound gives

$$\mathbb{E}_g \|g\|_{U^s}^{2^s} \leq \frac{\binom{2^s}{2}}{N}.$$

There is consequently a circle-valued g with at most this Gowers power. Put $\bar{g}, g$ at the two endpoints of the $(p + 1)$ -term pattern and put one at every intermediate position. The average is again one because $pd = 0$. Hence any factor valid for all such functions is at least

$$\left(\frac{N}{\binom{2^p}{2}} \right)^{1/2^p}.$$

The kernel product in the upper bound is exactly N , since the differences $1, \dots, p - 1$ are invertible. This proves the asserted optimal order. The random estimate is useful for large N ; for small N its right-hand side may exceed the trivial upper bound one. $\square$

A small explicit example also illustrates why the modulus restriction in the source theorem matters. On $\mathbb{Z}/3\mathbb{Z}$ take $g = (1, 1, -1)$ and the four functions $(\bar{g}, 1, 1, g)$. Their four-term average is one. However,

$$\|g\|_{U^3}^8 = \frac{49}{81} < 1.$$

Indeed, the zero derivative is one, and each nonzero derivative is a translate of $(1, -1, -1)$, whose U^2 fourth power is $(1/3)^4 + 2(2/3)^4 = 11/27$. The recursion gives $(1 + 2 \cdot 11/27)/3 = 49/81$.

[write] **6 October 2026; sources 06–10.** Rechecked at the second write: $(1/3)^4 + 2(2/3)^4 = 33/81 = 11/27$ and $(1 + 22/27)/3 = 49/81$. The $\mathbb{Z}/2\mathbb{Z}$ torsion examples of sources 03 and 04 (Part III above) are of the same kind; here the factor is quantified and shown sharp on $(\mathbb{Z}/2\mathbb{Z})^n$.

III.29 Density-sensitive counting and modular progression criteria

Source 10, delivered Section 6.

For convenience write

$$\Lambda_k(F_1, \dots, F_k) = \mathbb{E}_{x,d} \prod_{i=1}^k F_i(x + (i-1)d).$$

The following replacement for the counting step keeps the densities and uses only $k-2$ error terms.

Theorem III.29.1 (A density-weighted telescoping estimate). *Let $k \geq 3$ and let $A_1, \dots, A_k \subseteq G$. Put $F_i = \mathbf{1}_{A_i}$, $\delta_i = \mathbb{E}F_i$, and $b_i = F_i - \delta_i$. For $3 \leq j \leq k$ set*

$$\tau_j(G) = \left(\prod_{r=1}^{j-1} \kappa_G(r) \right)^{1/2^{j-1}}, \quad B_j = \left(\prod_{i=1}^{j-2} \delta_i^{1/2^i} \right) \delta_{j-1}^{1/2^{j-2}}, \quad D_j = \prod_{i=j+1}^k \delta_i.$$

Then

$$\left| \Lambda_k(F_1, \dots, F_k) - \prod_{i=1}^k \delta_i \right| \leq \sum_{j=3}^k D_j B_j \tau_j(G) \|b_j\|_{U^{j-1}}. \quad (\text{III.29.1})$$

Proof. For $L_i = x + (i-1)d$ use the pointwise identity

$$\prod_{i=1}^k F_i(L_i) - \prod_{i=1}^k \delta_i = \sum_{j=1}^k \left(\prod_{i < j} F_i(L_i) \right) b_j(L_j) \left(\prod_{i > j} \delta_i \right).$$

The averaged term with $j=1$ vanishes because $\mathbb{E}b_1 = 0$. The averaged term with $j=2$ also vanishes: the two adjacent forms are jointly uniform on G^2 , so its average factors and contains $\mathbb{E}b_2 = 0$. Apply Theorem III.28.2 in natural coefficient order to every remaining j -term average. The target is b_j , its kernel factor is $\tau_j(G)$, and $\|\mathbf{1}_{A_i}\|_{L^p} = \delta_i^{1/p}$ makes its mixed norm product exactly B_j . The triangle inequality proves the claim. $\square$

For one or two forms the count equals its product of densities exactly. If the exponent of G is coprime to $(k-1)!$, all the factors $\tau_j(G)$ are one. In particular, under the uniformity assumptions of Corollary 3.3 of the source paper, one has $\|b_j\|_{U^{j-1}} \leq \alpha$, and its error term can be replaced by

$$\alpha \sum_{j=3}^k D_j B_j \leq (k-2)\alpha.$$

The displayed density-dependent sum is often appreciably smaller.

Corollary III.29.2 (One-set counting). *Let $k \geq 3$, let $A \subseteq G$ have density δ , and put*

$$\eta = \|\mathbf{1}_A - \delta\|_{U^{k-1}}, \quad T_{k,G}(\delta) = \sum_{j=3}^k \delta^{k-j} \tau_j(G).$$

Then

$$|\Lambda_k(A) - \delta^k| \leq \delta \eta T_{k,G}(\delta). \quad (\text{III.29.2})$$

In the coprime case, put

$$S_k(\delta) = \sum_{m=0}^{k-3} \delta^m.$$

Then $T_{k,G}(\delta) = S_k(\delta)$ and

$$|\Lambda_k(A) - \delta^k| \leq \delta S_k(\delta) \eta. \quad (\text{III.29.3})$$

For $0 \leq \delta < 1$, $S_k(\delta) = (1 - \delta^{k-2})/(1 - \delta)$, while $S_k(1) = k - 2$.

Proof. The reciprocals in the definition of B_j sum to one, so $B_j = \delta$ when all the sets are A . Gowers norm monotonicity gives $\|\mathbf{1}_A - \delta\|_{U^{j-1}} \leq \eta$ for $j \leq k$. Theorem III.29.1 now gives both estimates.

For completeness, monotonicity follows from two applications of Cauchy–Schwarz. Writing $g = \Delta_{h_1} \cdots \Delta_{h_{s-1}} f$, one has

$$|\mathbb{E}g|^4 = \left| \mathbb{E}_t \mathbb{E}_x g(x) \overline{g(x+t)} \right|^2 \leq \mathbb{E}_t \left| \mathbb{E}_x g(x) \overline{g(x+t)} \right|^2 = \|g\|_{U^2}^4.$$

Average this inequality over the h_i , and use $(\mathbb{E}_h |\mathbb{E}g|^2)^2 \leq \mathbb{E}_h |\mathbb{E}g|^4$. The two sides identify with the appropriate cube forms and yield $\|f\|_{U^s} \leq \|f\|_{U^{s+1}}$. $\square$

Corollary III.29.3 (An explicit modular progression criterion). *Let $k \geq 3$, suppose that the exponent of G is coprime to $(k-1)!$, and let $A \subseteq G$ have positive density δ . With η as above, the number of ordered pairs (x, d) for which all k terms are distinct and lie in A is at least*

$$\delta N^2 \left(\delta^{k-1} - S_k(\delta) \eta - \frac{1}{N} \right). \quad (\text{III.29.4})$$

In particular, A contains such a progression if

$$N(\delta^{k-1} - S_k(\delta) \eta) > 1.$$

A simpler sufficient pair of conditions is

$$\eta \leq \frac{\delta^{k-1}}{2S_k(\delta)}, \quad N > 2\delta^{1-k}. \quad (\text{III.29.5})$$

Proof. There are exactly $|A| = \delta N$ counted pairs with $d = 0$. If $d \neq 0$ and two terms coincide, then $rd = 0$ for some $1 \leq r \leq k-1$, contradicting invertibility of multiplication by r . Subtracting these exactly δN degenerate pairs from (III.29.3), after multiplying by N^2 , proves (III.29.4). Either displayed condition makes this lower bound strictly positive. $\square$

In the notation of the source paper, being α -uniform of degree $k - 2$ means $\eta^{2^{k-1}} \leq \alpha$. Thus the sufficient uniformity condition in (III.29.5) can be written

$$\alpha \leq \left(\frac{\delta^{k-1}}{2S_k(\delta)} \right)^{2^{k-1}}, \quad N > 2\delta^{1-k}.$$

These improve the local modular-counting requirements. Incorporating them into a global quantitative Szemerédi argument also requires the later inverse, density-increment, and interval-transfer steps to be recomputed with their own hypotheses.

[write] **6 October 2026; sources 06–10.** Compared with the criteria of sources 03, 04 and 06 in the note after source 04’s Corollary III.4.4: the error here is δ times theirs, and in α terms the power of δ is $(k - 1)2^{k-1}$, between source 06’s $(k - \frac{1}{2})2^{k-2}$ and source 04’s $k2^{k-1}$. Not formalized.

III.29.1 Repeated terms and the appropriate main term

The torsion-sensitive estimate also gives a distinct-point criterion without a coprimality assumption. Define

$$\mathcal{D}_k(G) = \bigcup_{r=1}^{k-1} G[r].$$

A progression repeats a point exactly when $d \in \mathcal{D}_k(G)$. For each such difference, at most δN choices of x can contribute, since the first point must belong to A . Hence a sufficient criterion for a progression with k distinct points in an arbitrary finite abelian group is

$$\delta^{k-1} - \eta T_{k,G}(\delta) > \frac{|\mathcal{D}_k(G)|}{N}. \quad (\text{III.29.6})$$

In small characteristic it is sometimes preferable to reduce the pattern before estimating it. On $(\mathbb{Z}/p\mathbb{Z})^n$, if $k > p$, the homogeneous indicator product is exactly the p -term product, since coefficients repeat modulo p and indicators are idempotent. For a Bernoulli random set with inclusion probability δ , the exact expected normalized count is then

$$\left(1 - \frac{1}{N}\right) \delta^p + \frac{\delta}{N}.$$

Its leading independent-membership term is δ^p . This coefficient collision explains why a dimension-independent error tending to zero around δ^k cannot describe that pattern.

Further questions specific to these estimates. The sharpness argument determines the correct group-size order for $k = p + 1$; exact constants for larger k and general torsion profiles remain open within this investigation. A natural next problem is to retain the restricted images cG in Lemma III.28.1, instead of replacing each by all of G , and to exploit the overlap of these images. For heterogeneous sets, the telescoping order and the assignment of the L^p exponents can be optimized jointly; when all coefficient differences are invertible, the rearrangement inequality assigns the largest reciprocal exponents to the smallest densities within a fixed prefix. The finite fibre identity, mixed-norm induction, telescoping identity, and exact degenerate-pair count are also natural separate formalization targets.

III.30 An exact dimension threshold for polynomial colorings

Source 10, delivered Section 9.

There is a useful restricted question behind the recent constructions of sets with fewer arithmetic progressions than their uniformity might suggest: how many coordinates can a polynomial coloring control while forbidding reflection symmetry along every nontrivial progression? The construction of Shi and Dong uses a sum of odd-degree field norms [33]. We give an exact answer when the number of field-valued color coordinates and their total degrees are prescribed. The obstruction is a direct Chevalley–Warning argument [39, 40]; an analogous argument involving all homogeneous degrees appears in Fühner and Taranchuk [38]. Our purpose is to state the precise degree budget, prove its optimality, and identify the limit it places on this construction method. We make no claim of first discovery of these polynomial-method consequences.

Throughout this section, let $k = 2m \geq 4$ and let $p > k$ be prime. A coloring $Q : \mathbb{F}_p^d \rightarrow \mathbb{F}_p^r$ is *polynomial with coordinate degree bounds* $D_1, \dots, D_r$ if its ℓ th coordinate is represented by a polynomial Q_ℓ of total degree at most D_ℓ . A progression $(a, a+h, \dots, a+(k-1)h)$ is *symmetrically colored* if

$$Q(a+ih) = Q(a+(k-1-i)h) \quad (0 \leq i < m). \quad (\text{III.30.1})$$

It is nontrivial when $h \neq 0$. The assumption $p > k$ guarantees that all k of its points are then distinct.

III.30.1 Odd coefficients encode reflection symmetry

Fix a center $z \in \mathbb{F}_p^d$ and introduce formal variables $H = (H_1, \dots, H_d)$ and T . For a polynomial P of total degree at most D , write

$$P(z+TH) = \sum_{j=0}^D C_{j,z}(H)T^j. \quad (\text{III.30.2})$$

The polynomial $C_{j,z}$ is homogeneous of degree j , or is zero. Indeed, in the expansion of a monomial, selecting the coefficient of T^j selects exactly j factors from the coordinates of H . This coefficient definition needs no differentiation or division by factorials.

Lemma III.30.1 (Reflection and the odd coefficient equations). *Suppose $1 \leq D \leq k$ and put $s = \lceil D/2 \rceil$. For $z, h \in \mathbb{F}_p^d$, the following are equivalent:*

1. *The progression with center z and direction h is symmetrically colored by P .*
2. *$C_{1,z}(h) = C_{3,z}(h) = \dots = C_{2s-1,z}(h) = 0$.*
3. *The polynomial $P(z+Th)$ is even in T .*

Proof. Put $\alpha_i = i - \frac{1}{2} \in \mathbb{F}_p$ for $1 \leq i \leq m$. The progression with center z has points $z \pm \alpha_1 h, \dots, z \pm \alpha_m h$. Its reflection equalities are $P(z + \alpha_i h) = P(z - \alpha_i h)$.

The α_i are nonzero, and their squares are pairwise distinct: for $i \neq j$, the quantities $\alpha_i - \alpha_j = i - j$ and $\alpha_i + \alpha_j = i + j - 1$ are nonzero modulo p . In particular, $1 \leq i + j - 1 \leq k - 1 < p$.

Subtracting the two evaluations in (III.30.2) cancels all even powers, leaving

$$P(z+TH) - P(z-TH) = 2 \sum_{j=1}^s C_{2j-1,z}(H)T^{2j-1}.$$

Its degree is at most $2s - 1 \leq k - 1$, even when $D = k$. After division by $2\alpha_i$, the reflection equalities become

$$\sum_{j=1}^s C_{2j-1,z}(h)(\alpha_i^2)^{j-1} = 0 \quad (1 \leq i \leq m).$$

Since $s \leq m$, the first s equations form a square Vandermonde system with determinant

$$\prod_{1 \leq i < j \leq s} (\alpha_j^2 - \alpha_i^2) \neq 0.$$

Thus all its unknown odd coefficients vanish. Conversely, vanishing of these coefficients makes the polynomial even and implies every required reflection equality. $\square$

The important distinction is that the degree budget contains only $1, 3, \dots, 2s - 1$. An even top degree contributes no new reflection equation. This is why degree bounds $2s - 1$ and $2s$ lead to the same dimension threshold.

III.30.2 The zero-count obstruction

We recall the short finite-field proof of the form of Chevalley–Warning that we need.

Lemma III.30.2 (Chevalley–Warning in the needed form). *Let $F_1, \dots, F_t \in \mathbb{F}_p[H_1, \dots, H_d]$ vanish at the origin. Omit zero polynomials from the degree sum. If $\sum_{j=1}^t \deg F_j < d$, the number of their common zeros in $\mathbb{F}_p^d$ is divisible by p . Consequently there are at least p common zeros, including at least $p - 1$ nonzero common zeros.*

Proof. By Fermat’s theorem, the polynomial function

$$I(H) = \prod_{j=1}^t (1 - F_j(H)^{p-1})$$

equals the indicator of the common zero set. Its total degree is less than $d(p - 1)$. Every monomial of I therefore has at least one variable whose exponent e satisfies $0 \leq e < p - 1$. For that exponent,

$$\sum_{u \in \mathbb{F}_p} u^e = 0 \quad \text{in } \mathbb{F}_p.$$

For $e = 0$ this is $p = 0$; for $1 \leq e < p - 1$ it is the usual finite-field power sum. Summation of each monomial over $\mathbb{F}_p^d$ factors into the corresponding one-variable sums, and hence gives zero. It follows that $\sum_h I(h) = 0$ in $\mathbb{F}_p$, which is the required divisibility. The zero count is positive because the origin is a common zero, so it is at least p . $\square$

Theorem III.30.3 (Exact dimension threshold with heterogeneous degrees). *Let $r \geq 1$, let $1 \leq D_\ell \leq k$ for $1 \leq \ell \leq r$, and put*

$$s_\ell = \left\lceil \frac{D_\ell}{2} \right\rceil, \quad B = \sum_{\ell=1}^r s_\ell^2.$$

There exists a polynomial coloring $Q : \mathbb{F}_p^d \rightarrow \mathbb{F}_p^r$, with $\deg Q_\ell \leq D_\ell$, having no nontrivial symmetrically colored k -term arithmetic progression if and only if $d \leq B$.

If $d > B$, then for every such Q and every center z , the number of directions giving a symmetrically colored progression centered at z , including the zero direction, is divisible by p . In particular, there are at least $p - 1$ nonzero such directions at every center.

Proof. For necessity, fix z and apply Lemma III.30.1 to each color coordinate. Reflection symmetry is equivalent to the common vanishing of

$$C_{\ell,1,z}(h), C_{\ell,3,z}(h), \dots, C_{\ell,2s_\ell-1,z}(h) \quad (1 \leq \ell \leq r).$$

These homogeneous polynomials vanish at the origin and have total degree sum at most

$$\sum_{\ell=1}^r \sum_{j=1}^{s_\ell} (2j-1) = \sum_{\ell=1}^r s_\ell^2 = B.$$

If $d > B$, Lemma III.30.2 gives the asserted divisibility and nonzero directions. Thus avoidance requires $d \leq B$.

For sufficiency we use the odd-degree norm construction of Shi and Dong [33, Sections 3–4], allowing a different top degree in each coordinate. For a positive odd integer t , identify $\mathbb{F}_p^t$ linearly with $\mathbb{F}_{p^t}$ and let $N_t : \mathbb{F}_p^t \rightarrow \mathbb{F}_p$ be the field norm. In degree one, take $N_1(u) = u$. We need that N_t is homogeneous of degree t and that $N_t(u) = 0$ holds exactly when $u = 0$.

Here is a polynomial description of those properties. If $\beta_1, \dots, \beta_t$ is a basis of $\mathbb{F}_{p^t}$ over $\mathbb{F}_p$, then

$$N_t(u_1, \dots, u_t) = \prod_{a=0}^{t-1} \left(\sum_{j=1}^t u_j \beta_j^{p^a} \right).$$

This is a product of t linear forms, and Frobenius permutes its factors, so its coefficients lie in $\mathbb{F}_p$. Its value is the norm of $\sum_j u_j \beta_j$, which is nonzero precisely when that field element is nonzero. In particular, its formal total degree is t , despite the alternative large-exponent expression for a field norm.

For color coordinate ℓ , introduce disjoint variable blocks

$$u_\ell^{(1)} \in \mathbb{F}_p, \quad u_\ell^{(3)} \in \mathbb{F}_p^3, \quad \dots, \quad u_\ell^{(2s_\ell-1)} \in \mathbb{F}_p^{2s_\ell-1}.$$

There are s_ℓ^2 variables in that coordinate's collection of blocks. On their full B -dimensional product define

$$Q_\ell(u) = \sum_{j=1}^{s_\ell} N_{2j-1}(u_\ell^{(2j-1)}). \quad (\text{III.30.3})$$

Its total degree is $2s_\ell - 1 \leq D_\ell$.

Suppose a progression with center z and direction h is symmetrically colored. Lemma III.30.1 says that every $Q_\ell(z + Th)$ is even. If h has a nonzero block in coordinate ℓ , let t be the largest odd index with $h_\ell^{(t)} \neq 0$. All higher-index blocks contribute constants, because their directions vanish; all lower-index blocks have degree less than t . The coefficient of T^t in $Q_\ell(z + Th)$ is therefore exactly

$$N_t(h_\ell^{(t)}) \neq 0.$$

Since t is odd, this contradicts evenness. Thus every direction block vanishes, and $h = 0$. This proves avoidance at $d = B$.

For $d < B$, restrict the construction to any linear embedding of $\mathbb{F}_p^d$ into $\mathbb{F}_p^B$. A nonzero direction remains nonzero under the embedding, and composition with a linear map does not increase total degree, so all the required properties are preserved. $\square$

[write] **6 October 2026; sources 11–14.** Source 14 reproves this threshold and determines the exact minimum count above it, $q^d(q^{\max(d-B,0)} - 1)$ over every $\mathbb{F}_q$ with $p > k$, with subspace equality and a prime-field gap (its Theorem III.36.1); the quintic question remains open (item III.19).

[write] **7 October 2026; sources 60–69.** *Two later uses* (Part III, sources 64 and 67): source 64's dimension barrier $n \leq tk^2/4$ for degree $k-1$ (its Theorem III.107.2) is this theorem at $D_\ell = k-1$, uncited there; source 67 removes the degree cap for sums of independent norm blocks, whose dimension is at most rm^2 for $2m$ -term progressions (its Theorem III.109.1).

Corollary III.30.4 (Equal degree bounds and optimal coordinate budget). *For $1 \leq D \leq k$, a coloring by r polynomial coordinates of total degree at most D can avoid all nontrivial symmetrically colored k -APs on $\mathbb{F}_p^d$ exactly when*

$$d \leq r \lceil D/2 \rceil^2.$$

In particular, for degree at most $k = 2m$ the exact condition is $d \leq rm^2$, and for $d \geq 1$ the least possible number of polynomial coordinates is $\lceil d/m^2 \rceil$.

At the equality dimension $d = B$, the coloring (III.30.3) is surjective onto $\mathbb{F}_p^r$. Indeed, each coordinate contains its own independent linear variable $u_\ell^{(1)}$. After choosing all other variables, the r linear variables are uniquely determined by the desired color. Thus every color class has exactly p^{B-r} points.

III.30.3 An upper bound for arbitrary degree budgets

When the degree exceeds the progression length, the reflection equations can vanish without all odd coefficients vanishing. Nevertheless, they give a useful explicit dimension bound.

Proposition III.30.5 (A dimension bound in every degree). *For a positive integer s , define*

$$B_m(s) = \begin{cases} s^2, & s \leq m, \\ m(2s - m), & s \geq m. \end{cases}$$

Let $D_1, \dots, D_r \geq 1$ be arbitrary, put $s_\ell = \lceil D_\ell/2 \rceil$, and suppose $Q : \mathbb{F}_p^d \rightarrow \mathbb{F}_p^r$ is polynomial with $\deg Q_\ell \leq D_\ell$. If

$$d > \sum_{\ell=1}^r B_m(s_\ell),$$

then at every prescribed center there are at least $p-1$ nonzero directions producing symmetrically colored k -term progressions. Thus avoidance requires $d \leq \sum_{\ell=1}^r B_m(s_\ell)$.

Proof. Fix a center z and write the odd coefficients of one coordinate as $C_{2j+1,z}(H)$, for $0 \leq j < s$, with $s = \lceil D/2 \rceil$. Let $y_i = \alpha_i^2$ for the nonzero half-integer parameters used in Lemma III.30.1, and put $v_j = (y_i^j)_{i=1}^m \in \mathbb{F}_p^m$. After dividing each reflection equality by $2\alpha_i$, the vector of reflection equations is

$$E(H) = \sum_{j=0}^{s-1} v_j C_{2j+1,z}(H) = 0.$$

If $s \leq m$, these are equivalent to the individual odd coefficient equations, whose total degree sum is s^2 .

If $s > m$, form the square matrix

$$V = (v_{s-m}, v_{s-m+1}, \dots, v_{s-1}).$$

It is the usual Vandermonde matrix multiplied on the left by the invertible diagonal matrix with entries y_i^{s-m} , so it is invertible. The equivalent system $V^{-1}E(H) = 0$ has equations

$$F_j(H) = C_{2(s-m+j)+1,z}(H) + \sum_{q=0}^{s-m-1} b_{j,q} C_{2q+1,z}(H) = 0, \quad 0 \leq j < m,$$

for scalars $b_{j,q} \in \mathbb{F}_p$. Every coefficient in the second sum has strictly smaller degree than the displayed first coefficient. Consequently

$$\deg F_j \leq 2(s-m+j)+1, \quad \sum_{j=0}^{m-1} \deg F_j \leq m(2s-m).$$

Zero polynomials may again be omitted. Although the transformed equations need not be homogeneous, they all vanish at the origin. Applying this procedure independently to each color coordinate produces a system equivalent to reflection symmetry with total degree sum at most $\sum_{\ell} B_m(s_{\ell})$. The conclusion follows from Lemma III.30.2. $\square$

For four-term progressions and a scalar degree bound $D = 5$ or 6 , Proposition III.30.5 gives $d \leq 8$. The cubic construction already attains dimension 4, so the corresponding maximum dimension lies between 4 and 8. Determining that maximum is a concrete question beyond the exact low-degree theorem. No matching construction is asserted for the arbitrary-degree bound.

III.30.4 An explicit four-dimensional example

For $k = 4$, both scalar degree bounds $D = 3$ and $D = 4$ have exact dimension threshold 4. Over $\mathbb{F}_5$, the polynomial $T^3 + T + 1$ has no root and is therefore irreducible. If $\theta^3 + \theta + 1 = 0$, the norm of $a + b\theta + c\theta^2$ is

$$\begin{aligned} N_3(a, b, c) &= \det \begin{pmatrix} a & -c & -b \\ b & a-c & -b-c \\ c & b & a-c \end{pmatrix} \\ &= a^3 - 2a^2c + ac^2 + ab^2 + 3abc - b^3 - bc^2 + c^3 \pmod{5}. \end{aligned} \quad (\text{III.30.4})$$

The determinant is that of multiplication by $a + b\theta + c\theta^2$ in the basis $1, \theta, \theta^2$. Consequently

$$Q(x_0, x_1, x_2, x_3) = x_0 + N_3(x_1, x_2, x_3)$$

is a five-coloring of $\mathbb{F}_5^4$ with five color classes of 125 points and no nontrivial symmetrically colored 4-AP.

The accompanying exact-arithmetic verification checks the determinant formula, the absence of nonzero norm zeros, and all $625 \cdot 624 = 390,000$ ordered nontrivial progressions. It finds no symmetrically colored progression. This finite check verifies the displayed example; the theorem in arbitrary dimension and characteristic follows from the preceding proof.

III.30.5 What the threshold does and does not settle

The same result holds for any prescribed centrally symmetric pattern $\{\pm\alpha_1, \dots, \pm\alpha_m\}$ in $\mathbb{F}_p$ with $\alpha_i \neq 0$ and pairwise distinct α_i^2 . The conditions $Q(z + \alpha_i h) = Q(z - \alpha_i h)$ are again equivalent to the same odd coefficient equations, by the same Vandermonde argument. Hence the optimal dimension depends on the degree budgets and the number of reflection pairs, rather than on equal spacing itself.

For degree at most k , the theorem shows that the formal palette $\mathbb{F}_p^r$ must satisfy $r/d \geq 4/k^2$, and the equality construction uses all p^r colors. This is a sharp statement about a prescribed number of polynomial color coordinates. It is not a lower bound on the actual image size of arbitrary maps with redundant color coordinates, nor an optimality theorem for unrestricted colorings of intervals or cyclic groups. In particular, the auxiliary digit colors in Shi and Dong's cyclic construction are outside this polynomial coordinate model. The possibility of $N^{o(1)}$ colors in the Deng–Tidor–Zhao symmetric-progression problem remains outside the conclusion [37].

The restriction $D_\ell \leq k$ also matters. With higher odd degrees, agreement at the k progression parameters need not make the entire odd polynomial vanish. Thus higher-degree constructions and nonpolynomial color data provide specific directions beyond the class solved here. Within the stated class, however, both sides of the optimization problem are complete: the odd-degree sum gives a universal obstruction, and the known norm construction attains it.

III.31 Centered cube counts and their spectral decomposition

Source 13, delivered Section 7.

Lemma 3.10 bounds the normalized number of d -dimensional cubes in a set of density δ by $(\delta + \eta)^{2^d}$, where η is the appropriate uniformity norm of its balanced function. The triangle inequality proves that bound, but it treats all terms in the binomial expansion as possible. Mean zero forces many of them to vanish exactly.

In this section let $G = \mathbb{F}_p^n$, where p is an odd prime, and let $f : G \rightarrow \mathbb{R}$ have mean zero. Initially no absolute bound on f is required. Put $V_d = \{0, 1\}^d$, $m = 2^d$, and

$$L_v(x, h) = x + v \cdot h, \quad I_S(f) = \mathbb{E}_{x, h} \prod_{v \in S} f(L_v(x, h)), \quad I_\emptyset = 1.$$

For $\delta \geq 0$, define

$$\mathcal{Q}_d(\delta + f) = \mathbb{E}_{x, h} \prod_{v \in V_d} (\delta + f(L_v(x, h))) = \|\delta + f\|_{U^d}^m.$$

The last equality uses real-valuedness. For $f = \mathbf{1}_A - \delta$, this is the normalized count of all additive cubes, including degenerate ones, exactly as in the source lemma.

III.31.1 Cancellation from an affine coloop

Associate to each vertex v the column $\tilde{v} = (1, v) \in \mathbb{F}_p^{d+1}$. A vertex $v \in S$ is a *coloop* of S if its column is not in the span of the other columns in S . Call S *coloopless* if it has no coloop, with the empty set included by convention. This is the standard vector-matroid language for the affine dependencies of the cube; that matroid is a classical object [53]. The role of the terminology here is simply to identify a factor that can be averaged independently.

Lemma III.31.1 (Coloop cancellation). *If S has a coloop, then $I_S(f) = 0$. For every $S \subset V_d$,*

$$|I_S(f)| \leq \|f\|_{U^d}^{|S|}.$$

Proof. If v is a coloop, elementary linear algebra gives $w \in \mathbb{F}_p^{d+1}$ such that $\tilde{v} \cdot w = 1$ and $\tilde{u} \cdot w = 0$ for all other $u \in S$. For each $t \in G$, translate the $(d+1)$ variables by $(w_0 t, \dots, w_d t)$. Uniform measure is invariant under this translation. It changes only the factor at v , replacing it by $f(L_v + t)$. Average in t ; its average is zero. The norm bound follows from (I.25.2), placing f at vertices in S and the constant 1 elsewhere. $\square$

Let $c_{d,p}(j)$ count the coloopless j -subsets of V_d and write $\eta = \|f\|_{U^d}$. Direct expansion gives the exact formula

$$\mathcal{Q}_d(\delta + f) = \sum_{S \text{ coloopless}} \delta^{m-|S|} I_S(f). \quad (\text{III.31.1})$$

Thus the relevant coefficient polynomial is determined by the actual coefficient matroid over $\mathbb{F}_p$. For higher dimensions, one must not silently substitute characteristic-zero ranks.

III.31.2 The quartic term in every dimension

Lemma III.31.2 (Four-vertex classification). *Every three distinct cube vertices are affinely independent. In odd characteristic, an affinely dependent four-subset is exactly an ordinary parallelogram of the binary cube. Its expectation $I_S(f)$ is ρ^4 , where $\rho = \|f\|_{U^2}$.*

Proof. Reflect coordinates to move one vertex to 0. Two distinct nonzero binary vectors cannot be proportional: at a coordinate where the first is 1, the proportionality factor must be 1, forcing equality. This proves independence of three vertices.

For four dependent vertices, reflect one to 0 and call the others u, v, w . In a nonzero relation $\lambda u + \mu v + \nu w = 0$, each of $\lambda, \mu, \nu, \lambda + \mu + \nu$ is nonzero; otherwise three vertices would be dependent. A nonzero coordinate pattern among (u_i, v_i, w_i) can therefore only be 110, 101, or 011: singleton patterns and 111 are excluded. At least two pair patterns occur, or the vectors would coincide. Relabel so that 110 and 101 occur. They force $\mu = \nu = -\lambda$. Since $2\lambda \neq 0$, pattern 011 cannot occur. Consequently $u = v + w$ with disjoint supports for v, w . The four vertices form an ordinary parallelogram.

Any three forms belonging to that parallelogram are jointly uniform on G^3 ; their coefficient matrix has rank three. Hence

$$I_S(f) = \mathbb{E}_{a,b,c} f(a)f(b)f(c)f(a+b-c) = \sum_{\chi} |\hat{f}(\chi)|^4 = \rho^4.$$

$\square$

Proposition III.31.3 (The exact quartic coefficient). *The number of four-vertex parallelograms in V_d is*

$$P_d = \frac{6^d - 2 \cdot 4^d + 2^d}{8}. \quad (\text{III.31.2})$$

For $d \geq 2$,

$$\left| \mathcal{Q}_d(\delta + f) - \delta^m - P_d \delta^{m-4} \rho^4 \right| \leq \sum_{j=5}^m c_{d,p}(j) \delta^{m-j} \eta^j. \quad (\text{III.31.3})$$

An empty sum is zero; in particular $\mathcal{Q}_2(\delta + f) = \delta^4 + \rho^4$.

Proof. Coordinatewise, there are six ordered solutions to $a+b = c+e$ in $\{0, 1\}$. Thus the binary cube has 6^d ordered solutions. The degenerate ones, with the same two vertices on both sides, number $2 \cdot 4^d - 2^d$. Every nondegenerate parallelogram has eight orderings, proving (III.31.2). The values 0, 1, 2 remain distinct in odd characteristic. Equation (III.31.3) now follows from lemmas III.31.1 and III.31.2 and (III.31.1). $\square$

In particular, the error from the random-density main term starts at fourth order, rather than at first order. The exact quartic coefficient cannot be reduced for bounded real weights: replace f by tg for any fixed nonzero mean-zero real g , fix $0 < \delta < 1$, and let $t \downarrow 0$. For small t , the function $\delta + tg$ lies in $[0, 1]$. The quartic term is $P_d \delta^{m-4} t^4 \|g\|_{U^2}^4$, and all remaining terms have order at least t^5 . This is a weighted sharpness statement; it does not by itself assert an extremal theorem restricted to indicators at a fixed finite group size.

III.31.3 All surviving subsets of the three-cube

Proposition III.31.4 (The three-dimensional coefficient table). *For every odd p , the values $c_{3,p}(j)$, $j = 0, \dots, 8$, are*

$$(1, 0, 0, 0, 12, 8, 28, 8, 1).$$

Consequently, with $\eta = \|f\|_{U^3}$,

$$\mathcal{Q}_3(\delta + f) \leq \delta^8 + 12\delta^4\eta^4 + 8\delta^3\eta^5 + 28\delta^2\eta^6 + 8\delta\eta^7 + \eta^8. \quad (\text{III.31.4})$$

The quartic term may be replaced by $12\delta^4\rho^4$.

Proof. The first four entries and the twelve four-subsets follow from lemma III.31.2 and proposition III.31.3. Every affine hyperplane in $\mathbb{F}_p^3$ contains at most four binary vertices: choose a variable with a nonzero coefficient, assign the other two bits, and solve for that variable. Therefore any five vertices have full affine rank four after homogenization.

A five-subset has a coloop exactly when it contains a parallelogram. Each parallelogram extends to four five-subsets, and no five-subset contains two parallelograms. To justify the latter statement, two four-subsets of a five-set share three vertices, while three cube vertices complete to at most one parallelogram. After reflection, the possible completions of $0, u, v$ are $u+v, u-v, v-u$; at most one can be binary, according as the supports are disjoint or one strictly contains the other. Thus $12 \cdot 4 = 48$ of the $\binom{8}{5} = 56$ five-subsets have a coloop, leaving eight.

Deleting a vertex from a subset of size at least six still leaves at least five vertices and hence full rank. Every such subset is coloopless. Their counts are 28, 8, 1. Apply (III.31.1), lemma III.31.1, and $\rho \leq \eta$. $\square$

For comparison, the source bound $(\delta + \eta)^8$ has nonzero coefficients in the first three orders. Its quartic coefficient is 70, and its fifth coefficient is 56. Centering removes the first three nonconstant orders and reduces the next two coefficients to 12 and 8. These are exact combinatorial cancellations.

[write] **6 October 2026; sources 11–14.** *Credit.* The quartic coefficient P_d is source 03's R_d (Theorem III.2.6), source 04's P_d , source 08's R_d and source 09's P_d ; on odd-order groups the three-cube table $(1, 0, 0, 0, 12, 8, 28, 8, 1)$ is source 04's support count (Lemma III.3.3), source 08's Theorem III.9.2 and source 14's Appendix A (Section III.38). Here it is stated over $\mathbb{F}_p^n$; those sources prove it on every group of odd order (source 03 on every finite abelian group, with the order-two correction).

III.31.4 A complete spectral identity

Write $a_\xi = \widehat{f}(\xi)$, so $a_{-\xi} = \overline{a_\xi}$, and put

$$\sigma^2 = \mathbb{E}f^2 = \sum_{\xi} |a_\xi|^2, \quad \rho^4 = \sum_{\xi} |a_\xi|^4, \quad \zeta = \max_{\xi} |a_\xi|, \quad \eta = \|f\|_{U^3}.$$

Define the real quantities

$$T_5 = \sum_{\xi} |a_\xi|^2 a_\xi^2 \overline{a_{2\xi}}, \quad (\text{III.31.5})$$

$$B_1 = \sum_{\xi, \nu} |a_\xi|^2 |a_\nu|^2 |a_{\xi+\nu}|^2, \quad (\text{III.31.6})$$

$$B_2 = \sum_{\xi, \nu} |a_\xi|^2 a_\nu^2 a_{\xi-\nu} \overline{a_{\xi+\nu}}, \quad (\text{III.31.7})$$

$$B_3 = \sum_{\xi, \nu} a_\xi^2 a_\nu^2 a_{-\xi-\nu}^2. \quad (\text{III.31.8})$$

Finally let T_7 be the common expectation $I_S(f)$ of a seven-subset of the cube. Cube symmetries show that it is independent of the omitted vertex.

Theorem III.31.5 (Exact centered three-cube decomposition). *For every real mean-zero f on $\mathbb{F}_p^n$, p odd,*

$$\begin{aligned} \mathcal{Q}_3(\delta + f) &= \delta^8 + 12\delta^4 \rho^4 + 8\delta^3 T_5 \\ &\quad + \delta^2(12B_1 + 12B_2 + 4B_3) + 8\delta T_7 + \eta^8. \end{aligned} \quad (\text{III.31.9})$$

Moreover,

$$|T_5| \leq \zeta \rho^4, \quad |B_2|, |B_3| \leq B_1 \leq \eta^6, \quad |T_7| \leq \eta^4 \sqrt{B_1}. \quad (\text{III.31.10})$$

Proof. The eight five-circuits are the sets consisting of a vertex, its three graph neighbors, and its antipode. They form one symmetry orbit. Their unique affine relation is the sum of the three neighbors minus twice the vertex minus the antipode. Fourier expansion therefore gives (III.31.5) for every five-circuit.

The six-subsets split into three symmetry orbits, according to whether their omitted pair has Hamming distance one, two, or three. There are 12, 12, 4 such pairs, respectively. Fourier expansion imposes one equation for each of x, h_1, h_2, h_3 . The following kernel parametrizations, in lexicographic order of the retained vertices, verify the corresponding products:

Omitted pair	Six Fourier indices	Contribution
000, 001	$(u + v, -u - v, -u, u, -v, v)$	B_1
000, 011	$(u + v, -u - v, v, -u - 2v, u, v)$	B_2
000, 111	$(v, u, -u - v, -u - v, u, v)$	B_3

For the middle row set $\xi = u + v$, $\nu = v$. Each parametrization solves the four constraints and has two free coordinates; its integer pivot minor is a unit. Thus the calculation is valid over every odd field without a hidden divisor. The seven-subsets contribute $8\delta T_7$, and the full subset contributes η^8 . This proves (III.31.9).

The first bound in (III.31.10) follows termwise from (III.31.5). The bound $|B_3| \leq B_1$ also follows by taking absolute values. For B_2 , Cauchy–Schwarz bounds its absolute value by

$$\left(\sum_{\xi, \nu} |a_\xi|^2 |a_\nu|^2 |a_{\xi-\nu}|^2 \right)^{1/2} \left(\sum_{\xi, \nu} |a_\xi|^2 |a_\nu|^2 |a_{\xi+\nu}|^2 \right)^{1/2} = B_1,$$

using $|a_{-\nu}| = |a_\nu|$. Gowers–Cauchy–Schwarz applied to the corresponding six-subset gives $B_1 \leq \eta^6$.

To bound T_7 , set

$$A(h, k) = \mathbb{E}_x f(x+h) f(x+k) f(x+h+k),$$

$$B(h, k) = \mathbb{E}_x f(x) f(x+h) f(x+k) f(x+h+k).$$

Splitting the cube into parallel faces gives $T_7 = \mathbb{E}_{h,k} A(h, k) B(h, k)$. Expanding $\mathbb{E} A^2$ gives a cube missing an adjacent pair, so $\mathbb{E} A^2 = B_1$. Expanding $\mathbb{E} B^2$ gives the full cube, so $\mathbb{E} B^2 = \eta^8$. Cauchy–Schwarz proves the final bound. $\square$

[write] **6 October 2026; sources 11–14.** *Credit.* The identity (III.31.9) is source 08’s complete odd-order three-cube expansion (Theorem III.9.2), which holds on every finite abelian group of odd order; here it is proved over $\mathbb{F}_p^n$, p odd. New in this source: the orbit moments B_1, B_2, B_3, T_7 with the bounds (III.31.10), the spectral deficit below, and the sign examples.

III.31.5 A further sharp spectral deficit

The zero Fourier coefficient is absent because $\mathbb{E} f = 0$. Keeping this absence improves the elementary estimate $B_1 \leq \sigma^2 \rho^4$.

Proposition III.31.6 (A spectral upper bound with the zero-frequency loss). *If $\sigma^2 > 0$, then*

$$B_1 \leq \sigma^2 \rho^4 - \frac{\rho^8}{\sigma^2}. \quad (\text{III.31.11})$$

Equality is attained when $|\widehat{f}(\xi)|^2$ is constant for every nonzero ξ , including for a scalar multiple of $|G| \mathbf{1}_{\{0\}} - 1$.

Proof. Put $b_\xi = |a_\xi|^2$, so $b_0 = 0$ and $b_{-\xi} = b_\xi$. Translation gives the exact identity

$$\sigma^2 \rho^4 - B_1 = \frac{1}{2} \sum_{\nu} b_{\nu} \sum_{\xi} (b_{\xi} - b_{\xi+\nu})^2.$$

For each nonzero ν , retain only the two terms $\xi = 0, -\nu$. They contribute $2b_{\nu}^2$, whence the right side is at least $\sum_{\nu} b_{\nu}^3$. Cauchy–Schwarz yields

$$(\rho^4)^2 = \left(\sum_{\nu} b_{\nu}^2 \right)^2 \leq \left(\sum_{\nu} b_{\nu} \right) \left(\sum_{\nu} b_{\nu}^3 \right) = \sigma^2 \sum_{\nu} b_{\nu}^3.$$

This proves (III.31.11). If all nonzero b_{ν} are equal, the discarded translation terms vanish and the final Cauchy–Schwarz inequality is an equality. $\square$

For a balanced indicator, $\sigma^2 = \delta(1 - \delta)$ is known exactly. The right side of (III.31.11) is nonnegative, since $\rho^4 \leq \sigma^4$. Combining the proposition with theorem III.31.5 gives, for $\sigma > 0$,

$$B_1 \leq \mathcal{B} := \min \left\{ \eta^6, \sigma^2 \rho^4 - \rho^8 / \sigma^2 \right\}.$$

The case $\sigma = 0$ means $f = 0$ and is immediate.

One convenient two-sided consequence of the exact identity is

$$\begin{aligned} & \left| \mathcal{Q}_3(\delta + f) - \delta^8 - 12\delta^4 \rho^4 - 12\delta^2 B_1 - \eta^8 \right| \\ & \leq 8\delta^3 \zeta \rho^4 + 16\delta^2 B_1 + 8\delta \eta^4 \sqrt{B_1}. \end{aligned} \quad (\text{III.31.12})$$

If only an upper bound is wanted, theorem III.31.5 instead gives

$$\mathcal{Q}_3(\delta + f) \leq \delta^8 + 12\delta^4 \rho^4 + 8\delta^3 \zeta \rho^4 + 28\delta^2 \mathcal{B} + 8\delta \eta^4 \sqrt{\mathcal{B}} + \eta^8. \quad (\text{III.31.13})$$

This retains more information than the one-parameter polynomial (III.31.4); it is useful when the ordinary Fourier norm is much smaller than the higher uniformity norm.

III.31.6 Signs, characteristic, and the scope of sharpness

Odd-order circuit terms cannot simply be discarded. On $\mathbb{Z}/p\mathbb{Z}$, $p \geq 7$, take

$$g(x) = \cos(2\pi x/p) - u \cos(4\pi x/p), \quad u > 0.$$

Its only relevant Fourier modes are $\pm 1, \pm 2$, and (III.31.5) gives $T_5(g) = -u/16$. Replacing g by $-g$ reverses that sign. Multiplying by a sufficiently small number and adding a fixed $\delta \in (0, 1)$ makes this a valid $[0, 1]$ -valued weight. Likewise, for $g(x) = \cos(2\pi x/p) + u \sin(4\pi x/p)$, one has

$$B_3 = -3u^2/32 = -B_1.$$

Thus the possible negative contribution in (III.31.12) is real, even for elementary two-frequency examples.

The coloop cancellation mechanism is valid over every finite field, but the explicit odd-characteristic table is not. In characteristic two, the three-cube table is 1, 0, 0, 0, 14, 0, 28, 8, 1. Already its quartic coefficient differs. In larger dimensions one should compute $c_{d,p}(j)$ over the actual field and then apply (III.31.1). These rank questions connect the analytic problem to classical cube matroids [53], while control of more general linear systems belongs to the wider theory of true complexity [29].

Finally, neither the surviving-subset polynomial nor its spectral upper bound is asserted to be jointly optimal at every parameter tuple. We proved exact cancellation, exact identities, an exact leading coefficient, and explicitly stated inequalities with their own equality cases. Those are stronger and more precise claims than a numerical estimate, but they do not identify all simultaneous extremizers of the U^3 norm.

III.32 Interval discrepancy with the sharp length exponent

Source 13, delivered Section 8.

The proof of Lemma 3.5 bounds an interval's Fourier coefficient by frequency decay and explicitly discards the additional bound by the interval's length [1, 67]. Retaining both estimates gives a stronger local statement and an exact dual constant.

Theorem III.32.1 (Exact and explicit interval discrepancy bounds). *Let $N \geq 2$, $I \subset \mathbb{Z}/N\mathbb{Z}$ be an interval of length M , $0 \leq M \leq N$, and let f have mean zero. Put $\rho = \|f\|_{U^2}$, $\beta = M/N$, and $b = \min(\beta, 1 - \beta)$. Then*

$$|\mathbb{E}f\mathbf{1}_I| \leq \rho D_N(M), \quad D_N(M) = \left(\sum_{r=1}^{N-1} \left| \frac{\sin(\pi Mr/N)}{N \sin(\pi r/N)} \right|^{4/3} \right)^{3/4}. \quad (\text{III.32.1})$$

This constant is optimal over mean-zero real bounded functions. If $0 < M < N$, then

$$D_N(M) \leq \left(4b^{1/3} - 3N^{-1/3} \right)^{3/4} \leq 2\sqrt{2} b^{1/4}. \quad (\text{III.32.2})$$

For $M = 0$ or $M = N$, discrepancy is zero. The exponent $1/4$ cannot be replaced by a larger one, even for balanced indicators.

Proof. The geometric-series formula gives the absolute value of the normalized interval Fourier coefficient displayed in (III.32.1). Since $\widehat{f}(0) = 0$, Parseval and Hölder yield

$$|\mathbb{E}f\mathbf{1}_I| \leq \left(\sum_{r \neq 0} |\widehat{f}(r)|^4 \right)^{1/4} \left(\sum_{r \neq 0} |\widehat{\mathbf{1}}_I(r)|^{4/3} \right)^{3/4}.$$

This proves the exact upper bound. For its optimality, prescribe the Fourier coefficients of a test function by

$$\widehat{f}(r) = \widehat{\mathbf{1}}_I(r) |\widehat{\mathbf{1}}_I(r)|^{-2/3} \quad (r \neq 0), \quad \widehat{f}(0) = 0,$$

with zero used when the interval coefficient vanishes. These coefficients are conjugate symmetric, so f is real. They attain equality in Hölder. Scale f to have absolute value at most 1; the quotient of discrepancy by ρ is unchanged.

For the explicit estimate, nonzero coefficients are unchanged in absolute value when I is replaced by its complementary interval, so

$$|\widehat{\mathbf{1}}_I(r)| \leq \min\{b, (2|r|)^{-1}\},$$

where $|r|$ is the least positive centered frequency magnitude. The first bound is the interval length; the second follows from $|\sin(\pi r/N)| \geq 2|r|/N$. Let $F(x) = \min\{b, (2x)^{-1}\}^{4/3}$. This function is decreasing. Pair opposite frequencies, possibly overcounting the frequency $N/2$, and compare each summand with the integral on its preceding unit interval:

$$\sum_{r \neq 0} |\widehat{\mathbf{1}}_I(r)|^{4/3} \leq 2 \sum_{r=1}^{\lfloor N/2 \rfloor} F(r) \leq 2 \int_0^{N/2} F(x) dx.$$

For a nontrivial interval $b \geq 1/N$, the transition point $1/(2b)$ is at most $N/2$. Evaluation gives

$$2 \left[\frac{1}{2} b^{1/3} + 2^{-4/3} \int_{1/(2b)}^{N/2} x^{-4/3} dx \right] = 4b^{1/3} - 3N^{-1/3}.$$

Taking the $3/4$ power proves (III.32.2).

To prove the exponent is best possible, take $A = I$ with $\beta \rightarrow 0$ and $f = \mathbf{1}_I - \beta$. Then $\mathbb{E}f\mathbf{1}_I = \beta(1 - \beta)$, whereas Parseval and $|\widehat{\mathbf{1}}_I(r)| \leq \beta$ give $\rho^4 \leq \beta^3(1 - \beta)$. An estimate $|\mathbb{E}f\mathbf{1}_I| \leq C\rho\beta^\gamma$ with fixed C would therefore force $\gamma \leq 1/4$ as $\beta \downarrow 0$. $\square$

For $f = \mathbf{1}_A - \delta$ and the source convention of degree-one α -uniformity, $\rho \leq \alpha^{1/4}$. Multiplying the normalized estimate by N gives

$$||A \cap I| - \delta|I|| \leq 2\sqrt{2} \alpha^{1/4} N \min \left\{ \frac{|I|}{N}, 1 - \frac{|I|}{N} \right\}^{1/4}.$$

The repository's repaired Lemma 3.5 has the length-independent constant 3; the new expression improves it for every nontrivial interval and vanishes at both endpoint lengths. The exact dual constant is available if numerical constants matter. The same assertion holds for affine images of intervals under invertible slopes, because those maps permute Fourier frequencies. [6 October 2026; sources 11–14] The repaired Lemma 3.5 is the formal `lemma_3_5` (`Sections01_03.lean`, line 299; exact companion `lemma_3_5_holds`, `Proofs03Interval.lean`, line 412, at HEAD `e1db00e99`). No other source of this report treats Lemma 3.5.

Proposition III.32.2 (The small-relative-length limit). *If $M \rightarrow \infty$ and $M/N \rightarrow 0$, then*

$$\frac{D_N(M)}{(M/N)^{1/4}} \longrightarrow \left(\int_{\mathbb{R}} \left| \frac{\sin \pi t}{\pi t} \right|^{4/3} dt \right)^{3/4}, \quad (\text{III.32.3})$$

where the integrand takes its continuous value 1 at 0.

Proof. Write $\beta = M/N$ and rescale frequencies by $t = \beta r$. For bounded t , $r/N = t/M \rightarrow 0$, so the interval coefficient divided by β tends uniformly on compact sets to $\sin(\pi t)/(\pi t)$ in absolute value. The frequency mesh is $\beta \rightarrow 0$, giving the Riemann sum on every compact interval. The bound $(2|r|)^{-1}$ gives a tail, after division by $\beta^{1/3}$, of order $A^{-1/3}$ for $|r| > A/\beta$, uniformly in M, N . Let the compact cutoff A increase after taking the limit. The omission of frequency zero costs only β in the rescaled sum and disappears. Finally take the $3/4$ power. $\square$

This is the usual scaling analysis of a Dirichlet kernel, written here to identify the limiting optimal constant. The gain in the source lemma comes from retaining a parameter it had intentionally discarded.

III.33 Progression counting by telescoping

Source 13, delivered Section 9.

The 2^k factor in Corollary 3.3 comes from expanding every indicator into its mean and balanced part. A telescoping expansion needs only $k - 2$ errors, because the first two affine forms are exactly independent. We also record the precise group hypothesis needed by the argument.

Lemma III.33.1 (Endpoint generalized von Neumann inequality). *Let $k \geq 2$, and let G be a finite abelian group on which multiplication by each integer $1, \dots, k - 1$ is an automorphism. If $|f_i| \leq 1$, then*

$$\left| \mathbb{E}_{x,r \in G} \prod_{i=1}^k f_i(x + (i-1)r) \right| \leq \|f_k\|_{U^{k-1}}. \quad (\text{III.33.1})$$

The convention at $k = 2$ is $\|f\|_{U^1} = |\mathbb{E}f|$.

Proof. For $k = 2$, the map $(x, r) \mapsto (x, x + r)$ is bijective, and the average factors. Suppose the result holds at $k - 1$. Apply Cauchy–Schwarz in x and use $|f_1| \leq 1$. Expand the remaining square, write the two step variables as r and $r + u$, and replace $x + r$ by a new basepoint. With $\Delta_h f(y) = f(y)f(y + h)$, this gives

$$|\Lambda_k(f_1, \dots, f_k)|^2 \leq \mathbb{E}_u \left| \Lambda_{k-1}(\Delta_u f_2, \Delta_{2u} f_3, \dots, \Delta_{(k-1)u} f_k) \right|.$$

The induction hypothesis bounds the right side by $\mathbb{E}_u \left\| \Delta_{(k-1)u} f_k \right\|_{U^{k-2}}$. The power-mean inequality and bijectivity of multiplication by $k - 1$ bound this by

$$\left(\mathbb{E}_h \left\| \Delta_h f_k \right\|_{U^{k-2}}^{2^{k-2}} \right)^{1/2^{k-2}} = \|f_k\|_{U^{k-1}}^2.$$

The final equality is the defining derivative identity for Gowers norms. Take square roots. This is the normalized induction underlying Gowers’ Theorem 3.2 [1]. $\square$

The hypothesis holds on $\mathbb{F}_p^n$ whenever $p \geq k$, and on $\mathbb{Z}/N\mathbb{Z}$ whenever $\gcd(N, (k - 1)!) = 1$. Thus primality is a sufficient condition, not the underlying requirement. The condition cannot simply be dropped: on $\mathbb{Z}/3\mathbb{Z}$, take $k = 4$, $f_1 = f_4 = (1, 1, -1)$, and $f_2 = f_3 = 1$. The progression average is 1, since $x + 3r = x$ and $f_1(x)f_4(x) = 1$, whereas $\|f_4\|_{U^3}^8 = 49/81 < 1$. This gives an explicit failure of (III.33.1) without coefficient invertibility.

Theorem III.33.2 (A density-weighted telescoping bound). *Under the group hypothesis of lemma III.33.1, let $g_i : G \rightarrow [0, 1]$, $\delta_i = \mathbb{E}g_i$, and $f_i = g_i - \delta_i$. For $k \geq 3$,*

$$\left| \Lambda_k(g_1, \dots, g_k) - \prod_{i=1}^k \delta_i \right| \leq \sum_{j=3}^k \|f_j\|_{U^{j-1}} \prod_{i=j+1}^k \delta_i. \quad (\text{III.33.2})$$

In particular, if all the displayed norms are at most η , the error is at most $(k - 2)\eta$.

Proof. Replace g_k by $\delta_k + f_k$, then g_{k-1} by $\delta_{k-1} + f_{k-1}$ in the term containing δ_k , and continue. This is the exact identity

$$\Lambda_k(g_1, \dots, g_k) = \left(\prod_{i=3}^k \delta_i \right) \Lambda_2(g_1, g_2) + \sum_{j=3}^k \left(\prod_{i=j+1}^k \delta_i \right) \Lambda_j(g_1, \dots, g_{j-1}, f_j).$$

The first term equals $\prod_i \delta_i$ by two-form independence. Each balanced function is bounded by 1 in absolute value, so lemma III.33.1 bounds the j th error by $\|f_j\|_{U^{j-1}}$. Since each density is at most 1, there are at most $k - 2$ such errors. $\square$

The uniformity assumptions in the repository’s Corollary 3.3 say exactly that $\|f_j\|_{U^{j-1}} \leq \alpha$ for $j \geq 3$. Therefore (III.33.2) replaces its factor 2^k by

$$\sum_{j=3}^k \prod_{i=j+1}^k \delta_i \leq k - 2.$$

For unnormalized counts, multiply by $|G|^2$. [6 October 2026; sources 11–14] *Credit:* the generalized von Neumann bound and this telescope are source 03’s Lemma III.1.1 and Theorem III.1.2 and source 04’s Lemma III.4.2 and Theorem III.4.3, and also source 06’s Theorem III.5.2 and source 10’s

Theorem III.29.1; the corollary below is source 03's criterion (Corollary III.1.3; this S_k is its S_{k-2}). Sources 06 (quadratic in the norm) and 10 (error $\delta S_k \eta$) are stronger. Reversing the order gives a second valid bound with terms $\|f_j\|_{U^{k-j}} \prod_{i < j} \delta_i$, $1 \leq j \leq k-2$, but those norms require their own hypotheses; they are not supplied automatically by the forward assumptions.

Corollary III.33.3 (An explicit modular existence criterion). *Under the group hypothesis of lemma III.33.1, let $k \geq 3$, let $A \subset G$ have density $\delta > 0$, let $\eta = \|\mathbf{1}_A - \delta\|_{U^{k-1}}$, and set*

$$S_k(\delta) = \sum_{j=0}^{k-3} \delta^j.$$

If

$$\delta^k - \eta S_k(\delta) > \frac{\delta}{|G|}, \quad (\text{III.33.3})$$

then A contains a k -term progression with nonzero difference. Under the group hypothesis, its terms are distinct. A convenient sufficient pair of conditions is

$$\eta \leq \frac{\delta^k}{2S_k(\delta)}, \quad |G| > 2\delta^{1-k}.$$

Proof. Monotonicity of Gowers norms gives $\|\mathbf{1}_A - \delta\|_{U^{j-1}} \leq \eta$ for $3 \leq j \leq k$. Apply theorem III.33.2. Exactly $|A|$ pairs (x, r) have $r = 0$ and every term in A ; their normalized contribution is $\delta/|G|$. Inequality (III.33.3) leaves a nonzero-step pair. If two of its terms coincided, multiplication by an integer from 1 to $k-1$ would kill the nonzero difference, contradicting the automorphism hypothesis. $\square$

This is explicitly a modular/group conclusion. It does not, by itself, rectify a progression into the integer interval $[N]$. The source development has separate transfer modules for that purpose, and a global integer theorem must retain those steps.

III.34 The exact secondary term in the three-cube extremal problem

Source 14, delivered Section 2.

All averages are normalized. For a finite abelian group G and a real function $W : G \rightarrow \mathbb{R}$, put

$$Q_3(W) = \mathbb{E}_{x, h_1, h_2, h_3 \in G} \prod_{v \in \{0,1\}^3} W(x + v \cdot h) = \|W\|_{U^3}^8.$$

Throughout, $|G|$ is coprime to 6. The earlier leading-term result and its two-harmonic example left a gap: after subtracting $6\sqrt{2}\delta^4 u^4$, the coefficient at $\delta^{8/3} u^{16/3}$ was bounded below by $(9/4)^{1/3}$ and above by 6. We determine it exactly.

Theorem III.34.1 (Exact secondary coefficient, uniformly in the group). *Put $C_* = (9/4)^{1/3}$. There are absolute constants $c, C > 0$ such that, whenever $\delta > 0$, $f : G \rightarrow \mathbb{R}$ has mean zero, and $u = \|f\|_{U^3} \leq c\delta$,*

$$Q_3(\delta + f) \leq \delta^8 + 6\sqrt{2}\delta^4 u^4 + C_* \delta^{8/3} u^{16/3} + C \delta^{7/3} u^{17/3}. \quad (\text{III.34.1})$$

No pointwise bound on f is needed.

For $0 < \delta < 1$, define

$$M_\delta(u) = \sup_{\substack{G \text{ finite abelian, } (|G|,6)=1 \\ W:G \rightarrow [0,1], \mathbb{E}W=\delta \\ \|W-\delta\|_{U^3} \leq u}} Q_3(W).$$

Then, as $u \downarrow 0$,

$$M_\delta(u) = \delta^8 + 6\sqrt{2}\delta^4 u^4 + C_*\delta^{8/3}u^{16/3} + O_\delta(u^{17/3}). \quad (\text{III.34.2})$$

The upper constants are independent of G , even for varying groups.

This is a theorem about the local extremal cube problem. It does not by itself improve the final quantitative bound in Szemerédi’s theorem.

[write] **6 October 2026; sources 11–14.** *Source 04’s Question 1 is settled.* See Remark W4 at the end of this section for the deduction in source 04’s normalization. The “earlier project note” with the gap $[(9/4)^{1/3}, 6]$ is source 04 (Theorem III.3.8, Proposition III.3.9); its leading term $6\sqrt{2}$ and the norm gap (III.34.8) are source 04’s Theorems III.3.6 and III.3.4, reproved below; the centered expansion (III.34.7) is in sources 03, 04, 08, 09 and 13 (Appendix A, Section III.38, is its certificate here).

III.34.1 Fourier preliminaries and the previous leading estimate

Write the dual group additively, and put

$$a_r = \widehat{f}(r) = \mathbb{E}_x f(x) \overline{r(x)}, \quad a_{-r} = \overline{a_r}, \quad a_0 = 0.$$

We use the exact identities

$$\|f\|_{U^2}^4 = \sum_r |a_r|^4 =: S, \quad (\text{III.34.3})$$

$$\|f\|_{U^3}^8 = \sum_{s,t} |T(s,t)|^2, \quad T(s,t) = \sum_r a_r \overline{a_{r+s}} \overline{a_{r+t}} a_{r+s+t}, \quad (\text{III.34.4})$$

$$B_5(f) := \mathbb{E}_{x,h_1,h_2,h_3} f(x)f(x+h_1)f(x+h_2)f(x+h_3)f(x+h_1+h_2+h_3) \quad (\text{III.34.5})$$

$$= \sum_r a_{2r} a_{-r}^3 a_r. \quad (\text{III.34.6})$$

The sum in (III.34.4) is unnormalized. On inserting Fourier inversion in the cube, the four orthogonality equations say that the parity-adjusted frequency labels form an affine three-cube; grouping its two parallel faces gives the identity.

The centered expansion is

$$Q_3(\delta + f) = \delta^8 + 12\delta^4 S + 8\delta^3 B_5(f) + R_6 + R_7 + u^8, \quad |R_6| \leq 28\delta^2 u^6, \quad |R_7| \leq 8\delta u^7. \quad (\text{III.34.7})$$

Indeed, one, two, or three distinct cube forms are jointly uniform, so their centered terms vanish. The twelve four-element parallelograms contribute S each; all other four-element supports are independent. The eight surviving five-element supports are cube-symmetric copies of $\{000, 100, 010, 001, 111\}$ and give $B_5(f)$. Bound the 28 sextic and 8 septic supports by cube Cauchy–Schwarz. Odd order suffices: the nonzero minors in these support arguments are ± 1 or ± 2 .

Let $\lambda_r = |a_r|^2$ and $M = \sum_r |a_r|^8$. The coordinate axes in (III.34.4) contribute

$$2\mathcal{E}(\lambda) - S^2, \quad \mathcal{E}(\lambda) = \sum_{r+s=t+w} \lambda_r \lambda_s \lambda_t \lambda_w.$$

Retaining the union of

$$(r = t, s = w), \quad (r = w, s = t), \quad (s = -r, w = -t)$$

in this positive energy sum gives $3S^2 - 3M$. Each family weighs S^2 , each pairwise intersection weighs M , and the triple intersection has $2r = 0$ and zero weight. Consequently,

$$\|f\|_{U^3}^8 \geq 5S^2 - 6M \geq 2S^2. \quad (\text{III.34.8})$$

The final inequality uses conjugate pairing, which gives $M \leq S^2/2$. Equality for nonzero f forces exactly one conjugate pair.

Choose a largest pair $\{\xi, -\xi\}$, and write

$$\rho = |a_\xi|, \quad D = \sum_{r \notin \{\xi, -\xi\}} |a_r|^4, \quad b = |a_{2\xi}|, \quad S = 2\rho^4 + D. \quad (\text{III.34.9})$$

There is no order-three character, so the pairs $\{\pm\xi\}$ and $\{\pm 2\xi\}$ are disjoint, and $2b^4 \leq D$. The axes estimate yields

$$\text{axes contribution} \geq 8\rho^8 + 20\rho^4 D + 2D^2 = 2S^2 + 12\rho^4 D, \quad (\text{III.34.10})$$

because the tail eighth-power mass is at most $D^2/2$. Maximality of ρ also gives

$$\|f\|_{U^3}^8 \geq (2 + 3D/S)S^2, \quad (\text{III.34.11})$$

by using $M \leq \rho^4 S = (S - D)S/2$ in (III.34.8).

III.34.2 Eight additional Fourier slots

Lemma III.34.2 (The off-axis correction). *With the notation above,*

$$\|f\|_{U^3}^8 \geq 2S^2 + 12\rho^4 D + 8\rho^4 b^4 - 16\rho^3 D^{5/4} - 16\rho^2 D^{3/2}. \quad (\text{III.34.12})$$

Proof. Besides the axes, retain the eight slots

$$(s, t) \in \{(\pm\xi, \pm 3\xi), (\pm 3\xi, \pm\xi)\}.$$

They are distinct and off the axes: $3\xi \neq 0$ and $\xi \neq \pm 3\xi$. Their T -values have the same modulus, by $T(s, t) = T(t, s)$ and $T(-s, t) = \overline{T(s, t)}$.

Let $a'_r = a_r$ off $\{\xi, -\xi\}$ and $a'_r = 0$ on that pair. There is an exact decomposition

$$T(\xi, 3\xi) = \rho^2 b^2 + 2 \operatorname{Re}(a_\xi \overline{a_{2\xi}} \overline{a_{4\xi}} a_{5\xi}) + T_{a'}(\xi, 3\xi). \quad (\text{III.34.13})$$

The base $r = -2\xi$ gives labels $(-2\xi, -\xi, \xi, 2\xi)$ and contributes $\rho^2 b^2$. The only other nonzero summands involving a dominant label have bases $r = \xi$ and $r = -5\xi$; they are conjugates. The other possible bases force a zero label.

To check small torsion explicitly, the four offsets are 0, 1, 3, 4. Two dominant labels differ by $\pm 2\xi$. Offset difference 2 gives the displayed main term. The only extra coincidence in an order coprime to 6 is order 5, where offset difference 3 equals -2 ; both extra summands contain a_0 . For order 5 the middle term in (III.34.13) vanishes as well, since $5\xi = 0$. For all other allowed orders, its three other labels are in the tail.

Each tail coefficient is at most $(D/2)^{1/4}$. Hölder gives $|T_{a'}(\xi, 3\xi)| \leq D$. Thus

$$|T(\xi, 3\xi) - \rho^2 b^2| \leq 2\rho(D/2)^{3/4} + D.$$

Using $|A + E|^2 \geq A^2 - 2A|E|$ for $A \geq 0$,

$$\begin{aligned} 8|T(\xi, 3\xi)|^2 &\geq 8\rho^4 b^4 - 32\rho^3 b^2 (D/2)^{3/4} - 16\rho^2 b^2 D \\ &\geq 8\rho^4 b^4 - 16\rho^3 D^{5/4} - 16\rho^2 D^{3/2}. \end{aligned}$$

The last line uses $b^2 \leq (D/2)^{1/2}$ and rounds constants upward. Add this to (III.34.10). $\square$

If the tail consists of the doubled pair alone, $D = 2b^4$. The axes penalty above $2S^2$ is then $24\rho^4 b^4$; the eight additional terms raise it to $32\rho^4 b^4$, the exact leading penalty in the two-harmonic model.

III.34.3 Uniform localization and optimization

Proof of the upper estimate. The zero case is immediate. Scale to $g = f/u$, $\|g\|_{U^3} = 1$, and $t = u/\delta$, so that $Q_3(\delta + f) = \delta^8 Q_3(1 + tg)$. In this proof the Fourier variables refer to g . Put

$$S_0 = 1/\sqrt{2}, \quad \eta = S_0 - S \geq 0, \quad F = Q_3(1 + tg) - 1 - 6\sqrt{2}t^4.$$

For $0 < t \leq 1$, the centered expansion gives

$$F = -12\eta t^4 + 8t^5 B_5(g) + R, \quad |R| \leq 37t^6. \quad (\text{III.34.14})$$

If $F \leq 0$, the desired upper bound is immediate. Assume $F > 0$.

First $|B_5(g)| \leq \rho S < 1$. Hence $\eta = O(t)$ uniformly, and $S \geq S_0/2$ for sufficiently small absolute t . By (III.34.11),

$$3DS \leq 1 - 2S^2 = 2\eta(S_0 + S) \leq 2\sqrt{2}\eta, \quad D \leq (8/3)\eta.$$

The selected pair in B_5 contributes at most $2\rho^4(D/2)^{1/4}$, while the other terms contribute at most ρD . Since $2\rho^4 \leq S \leq S_0 < 1$, $D < 1$, and $\rho < 1$, each of these two contributions is at most $D^{1/4}$. Thus $|B_5(g)| \leq 2D^{1/4}$. Therefore

$$12\eta \leq 16tD^{1/4} + 37t^2 \leq 16(8/3)^{1/4}t\eta^{1/4} + 37t^2.$$

If $\eta \leq (37/6)t^2$ the desired stronger localization follows already. Otherwise subtract $37t^2 \leq 6\eta$ and divide by $\eta^{1/4}$. Both cases give, for example,

$$\eta \leq 7t^{4/3}, \quad D \leq 19t^{4/3}, \quad (\text{III.34.15})$$

for all sufficiently small absolute t .

Substitute $\rho^4 = (S_0 - \eta - D)/2$ into Lemma III.34.2, and use $2b^4 \leq D$:

$$\begin{aligned} 2\sqrt{2}\eta &\geq 3\sqrt{2}D + 2\sqrt{2}b^4 - (6D + 4b^4)(\eta + D) \\ &\quad - 16\rho^3 D^{5/4} - 16\rho^2 D^{3/2}. \end{aligned}$$

Here the nonnegative term $2\eta^2$ was discarded. By (III.34.15), all errors are uniformly $O(t^{5/3})$, so

$$\eta \geq \frac{3}{2}D + b^4 - O(t^{5/3}). \quad (\text{III.34.16})$$

If $a_\xi = \rho e^{i\theta}$, the selected quintic term is exactly $2\rho^4 \operatorname{Re}(a_{2\xi} e^{-2i\theta})$. Thus

$$B_5(g) \leq 2\rho^4 b + \rho D \leq S_0 b + D.$$

Consequently,

$$\begin{aligned} F &\leq -18Dt^4 - 12b^4t^4 + 4\sqrt{2}bt^5 + O(t^{17/3}) \\ &\leq -48b^4t^4 + 4\sqrt{2}bt^5 + O(t^{17/3}). \end{aligned}$$

The error scales are $t^4t^{5/3} = t^{17/3}$, $t^5D = O(t^{19/3})$, and $t^6 = O(t^{17/3})$. Finally,

$$\sup_{b \geq 0} (4\sqrt{2}tb - 48b^4) = C_* t^{4/3}, \quad b = \beta_* t^{1/3}, \quad \beta_* = \left(\frac{\sqrt{2}}{48} \right)^{1/3}.$$

Indeed $b^3 = \sqrt{2}t/48$, and $3\sqrt{2}\beta_* = (9/4)^{1/3}$. Rescaling proves (III.34.1), including uniformity in G . $\square$

III.34.4 The matching model and the supremum

On a character of order greater than 8, take $a_{\pm\xi} = \rho$, $a_{\pm 2\xi} = b$, all positive, and no others. There is no aliasing in frequency parallelograms, whose two-label sums lie in $[-4, 4]$. Here is an explicit enumeration certificate. In the Fourier cube identity, every frequency belongs to $\{-2, -1, 0, 1, 2\}$; the absence of aliasing means that all the face parallelograms close over the integers. Hence the frequency labels are integer affine cubes in this five-point interval.

Reflect each direction with negative increment. This makes all three increments nonnegative without changing the numbers of vertices of each absolute frequency. Their sum is at most 4. The possible multisets of positive increments are

$$\emptyset, (1), (2), (3), (4), (1, 1), (1, 2), (1, 3), (2, 2), (1, 1, 1), (1, 1, 2).$$

For a multiset of k positive increments with sum H and multiplicities $m_1, m_2, \dots$, choose the coordinate positions and the original signs in

$$\frac{3!}{(3-k)! \prod_j m_j!} 2^k$$

ways. The minimum vertex ranges over the integers $-2, \dots, 2-H$. This gives the following complete table, where n_j counts the vertices of absolute frequency j :

(n_0, n_1, n_2)	number of frequency cubes
$(8, 0, 0)$	1
$(4, 4, 0)$	24
$(4, 0, 4)$	24
$(3, 4, 1)$	16
$(2, 4, 2)$	96
$(0, 8, 0)$	8
$(0, 4, 4)$	48
$(0, 0, 8)$	8

Each row contributes its count times $\delta^{n_0} \rho^{n_1} b^{n_2}$. Consequently,

$$\|f\|_{U^3}^8 = 8\rho^8 + 48\rho^4 b^4 + 8b^8, \quad (\text{III.34.17})$$

$$Q_3(\delta + f) = \delta^8 + 24\delta^4(\rho^4 + b^4) + 16\delta^3\rho^4 b + 96\delta^2\rho^4 b^2 + \|f\|_{U^3}^8. \quad (\text{III.34.18})$$

These are the earlier cosine identities with amplitudes $2\rho, 2b$. Choose

$$b = \beta_* \delta^{-1/3} u^{4/3}, \quad \rho^4 = -3b^4 + \sqrt{u^8/8 + 8b^8}.$$

Then $\|f\|_{U^3} = u$ exactly and $\rho^4 = u^4/(2\sqrt{2}) - 3b^4 + O(b^8/u^4)$. Substitution gives

$$Q_3(\delta + f) = \delta^8 + 6\sqrt{2}\delta^4 u^4 + C_* \delta^{8/3} u^{16/3} + O_\delta(u^{20/3}).$$

For fixed $0 < \delta < 1$ these are $[0, 1]$ weights for small u . A fixed prime cyclic group of order 17 supplies them. This proves the matching lower bound in (III.34.2). For norms smaller than the stated budget u , apply the upper estimate at the actual norm and use monotonicity of its positive displayed powers.

III.34.5 Rigidity at the secondary scale

Theorem III.34.3 (Two-harmonic rigidity). *Fix $\delta > 0$. Let G_j have order coprime to 6 and let f_j be real and centered, with $u_j = \|f_j\|_{U^3} \downarrow 0$. Suppose*

$$Q_3(\delta + f_j) = \delta^8 + 6\sqrt{2}\delta^4 u_j^4 + C_* \delta^{8/3} u_j^{16/3} + o_\delta(u_j^{16/3}).$$

Then largest pairs $\{\pm \xi_j\}$, with $a_{j,\xi_j} = \rho_j e^{i\theta_j}$, satisfy

$$|a_{j,2\xi_j}| \sim \beta_* \delta^{-1/3} u_j^{4/3}, \quad (\text{III.34.19})$$

$$\frac{a_{j,2\xi_j}}{|a_{j,2\xi_j}|} e^{-2i\theta_j} \rightarrow 1, \quad (\text{III.34.20})$$

$$\sum_{r \notin \{\pm \xi_j, \pm 2\xi_j\}} |a_{j,r}|^4 = o_\delta(u_j^{16/3}), \quad (\text{III.34.21})$$

$$\rho_j^4 = \frac{u_j^4}{2\sqrt{2}} - 3\beta_*^4 \delta^{-4/3} u_j^{16/3} + o_\delta(u_j^{16/3}). \quad (\text{III.34.22})$$

After removing the two conjugate pairs, the remainder is $o_\delta(u_j^{4/3})$ in U^2 . No L^2 approximation is asserted.

Proof. Use the normalized variables above. Put $\beta = b/t^{1/3}$ and $e = (D - 2b^4)/t^{4/3} \geq 0$. Retaining both the phase and the unused tail in the upper proof gives

$$\frac{F}{t^{16/3}} \leq 4\sqrt{2}\beta \cos \phi - 48\beta^4 - 18e + o(1), \quad \phi = \arg a_{2\xi} - 2 \arg a_\xi.$$

Here replacing $2\rho^4 b \cos \phi$ by $S_0 b \cos \phi$ is an approximation with absolute error at most $(\eta + D)b$, because $2\rho^4 = S_0 - \eta - D$. We use this error estimate, rather than a one-sided inequality whose direction could depend on $\cos \phi$. After the t^5 factor and division by $t^{16/3}$, its contribution is $O(t^{4/3}) = o(1)$ by localization. Localization bounds β, e . The polynomial $4\sqrt{2}\beta - 48\beta^4$ has unique

maximum C_* at $\beta_* > 0$. Asymptotic equality forces $\beta \rightarrow \beta_*$, $e \rightarrow 0$, and $\cos \phi \rightarrow 1$. The exact centered expansion then gives

$$\eta/t^{4/3} \rightarrow (4\sqrt{2}\beta_* - C_*)/12 = 4\beta_*^4, \quad D/t^{4/3} \rightarrow 2\beta_*^4.$$

Use $\rho^4 = (S_0 - \eta - D)/2$ and rescale. The Fourier formula for U^2 gives the claimed remainder norm. $\square$

III.34.6 Indicator sharpness

Proposition III.34.4. *For every $0 < \delta < 1$, there are prime cyclic groups and subsets A_j with densities $\delta_j \rightarrow \delta$ and $v_j = \|1_{A_j} - \delta_j\|_{U^3} \downarrow 0$, such that*

$$\frac{Q_3(1_{A_j}) - \delta_j^8 - 6\sqrt{2}\delta_j^4 v_j^4}{\delta_j^{8/3} v_j^{16/3}} \rightarrow C_*.$$

Proof. Take the two-harmonic weights at $u_j = 2^{-j}u_0$, for a sufficiently small fixed $u_0 > 0$, on sufficiently large primes p_j . Their exact formulas do not depend on the prime once it exceeds 8. Include points independently with the prescribed weights. At most $28/p_j$ of normalized three-cubes have a repeated vertex. Elsewhere independence gives the weighted expectation exactly, for both the indicator cube polynomial and the product centered at the fixed mean δ . Expectation errors are therefore $O(1/p_j)$. Changing one inclusion decision changes either polynomial by at most $8/p_j$, so a bounded-difference variance estimate gives $O(1/p_j)$ variance. The density has variance at most $1/(4p_j)$. Therefore, for large enough p_j , some realization approximates all three statistics simultaneously to any prescribed error. Recentering at the actual density costs at most $8|\delta_j - \delta|$.

Take errors $o(u_j^{28/3})$. Then $v_j^8 = u_j^8 + o(u_j^{28/3})$, whence $v_j^4 = u_j^4 + o(u_j^{16/3})$. The model expansion transfers with its exact secondary coefficient, including the replacements of δ by δ_j . Since $v_j/u_j \rightarrow 1$ and $u_{j+1}/u_j = 1/2$, the actual norms v_j decrease strictly after discarding finitely many initial indices. $\square$

Remark III.34.5 (Why boundedness alone cannot give L^2 rigidity). The indicator examples also supply an explicit obstruction to replacing the U^2 remainder in Theorem III.34.3 by an L^2 remainder tending to zero. For those examples,

$$\|1_{A_j} - \delta_j\|_2^2 = \delta_j(1 - \delta_j) \rightarrow \delta(1 - \delta) > 0.$$

Let P_j be any Fourier projection onto at most four frequencies, including a projection onto two conjugate pairs. Each Fourier coefficient has modulus at most $v_j = \|1_{A_j} - \delta_j\|_{U^3}$, by (III.34.8). Thus Parseval gives $\|P_j\|_2^2 \leq 4v_j^2 \rightarrow 0$ and, by orthogonality,

$$\|1_{A_j} - \delta_j - P_j\|_2^2 \rightarrow \delta(1 - \delta).$$

Thus even exact asymptotic secondary extremality among bounded weights does not imply small L^2 distance from the two-harmonic projection.

Remark W4 (Source 04's Question 1; added at the third write, 6 October 2026). *Claim.* In source 04's normalization the limit in its display (Q1) exists and equals $C_* = (9/4)^{1/3}$: for $0 < \delta < 1$, with

$M_\delta^{04}(u)$ the supremum of $Q_3(W)$ over prime cyclic groups of order greater than 16 and weights $0 \leq W \leq 1$ with $\mathbb{E}W = \delta$ and $\|W - \delta\|_{U^3} = u$,

$$\lim_{u \downarrow 0} \frac{M_\delta^{04}(u) - \delta^8 - 6\sqrt{2}\delta^4 u^4}{\delta^{8/3} u^{16/3}} = (9/4)^{1/3}.$$

Proof. Every prime $p > 16$ is coprime to 6, and $\|W - \delta\|_{U^3} = u$ implies $\leq u$, so $M_\delta^{04}(u) \leq M_\delta(u)$, the supremum of Theorem III.34.1; by (III.34.2) the limsup of the displayed quotient is at most C_* . Source 04's Proposition III.3.9 gives, on $\mathbb{Z}/p\mathbb{Z}$ with $p > 16$ and norm exactly u , the quotient $(9/4)^{1/3} + o(1)$, so the liminf is at least $(9/4)^{1/3} = C_*$. *The normalizations agree:* source 04's model $f = a \cos(2\pi x/p) + b \cos(4\pi x/p)$ has Fourier coefficients $a/2$ and $b/2$, so this source's ρ, b are source 04's $a/2, b/2$; substituting $\rho = a/2$, $b = b_{04}/2$ in (III.34.17)–(III.34.18) gives source 04's identity (III.3.24) term by term, and $\beta_* = (\sqrt{2}/48)^{1/3}$ is half of source 04's $c = (\sqrt{2}/6)^{1/3}$. (Checked at the write: the model identity against a direct cube sum on $\mathbb{Z}/17\mathbb{Z}$ at two parameter points, agreement to 10^{-15} ; the normalized quotient along the model at $\delta = 1/4$ is 1.3104628 at $u = 10^{-4}$ and 1.3103709 at $u = 10^{-6}$, against $C_* = 1.3103707$.) This is a new result of source 14, not a second route; it is unrefereed and not formalized. The constants c, C of the upper bound are not explicit (source 14's Question 3). The concentration question of source 04's Question 2 is answered in U^2 by Theorem III.34.3 and fails in L^2 by the remark above.

III.35 Order five changes the next exponent

Source 14, delivered Section 3.

Even after excluding order-three characters, short additive resonances still matter. In particular, order five produces a further distinction: on groups of exponent five, the correction of order $u^{16/3}$ vanishes, and the next exponent is 6. The optimal coefficient at u^6 is not determined here.

Lemma III.35.1 (An order-five energy penalty). *If the largest Fourier character ξ has order 5, then*

$$\|f\|_{U^3}^8 \geq 2S^2 + 16\rho^6 b^2. \quad (\text{III.35.1})$$

Proof. In addition to the three energy families retained in the norm-gap proof, retain the additive quadruple

$$(\xi, \xi, -\xi, -2\xi).$$

It is additive because $2\xi = -3\xi$. Interchanging the sides, interchanging the unequal entries on a side, and negating all entries produces eight distinct quadruples, all outside the previous three families. Each has weight $\rho^6 b^2$. Thus $\mathcal{E}(\lambda) \geq 3S^2 - 3M + 8\rho^6 b^2$. The axes bound and $M \leq S^2/2$ prove the claim. $\square$

Theorem III.35.2 (Sharp correction exponent in exponent five). *Let $\delta > 0$, and let G be a nontrivial finite abelian group of exponent 5. Uniformly over such groups and real centered functions,*

$$Q_3(\delta + f) \leq \delta^8 + 6\sqrt{2}\delta^4 u^4 + O(\delta^2 u^6) \quad (u/\delta \downarrow 0).$$

For $0 < \delta < 1$, the exponent 6 cannot be increased, even on the fixed group $\mathbb{Z}/5\mathbb{Z}$. This group has $[0, 1]$ weights of mean δ , with centered norm u , satisfying

$$Q_3(W_u) = \delta^8 + 6\sqrt{2}\delta^4 u^4 + \frac{2^{3/4}}{3} \delta^2 u^6 + O_\delta(u^8). \quad (\text{III.35.2})$$

The coefficient $2^{3/4}/3$ is a lower bound, not a claim of optimality.

Uniform upper estimate. Normalize by $g = f/u$, $t = u/\delta$, as in the preceding proof, and use F, η, D, ρ, b there. If $F \leq 0$ there is nothing to prove. Otherwise the first localization step gives $\eta = O(t)$ and $D = O(\eta)$. Thus $\rho^4 = (S_0 - \eta - D)/2$ is bounded below by a positive absolute constant. By Lemma III.35.1, $b^2 = O(\eta)$, since $1 - 2S^2 \leq 2\sqrt{2}\eta$. The quintic estimate becomes

$$|B_5(g)| \leq 2\rho^4 b + \rho D = O(\eta^{1/2} + \eta).$$

The centered cube expansion gives

$$\eta \leq C_1 t \eta^{1/2} + C_2 t \eta + C_3 t^2.$$

For small t , absorb the middle term. The resulting quadratic inequality in $\sqrt{\eta}$ gives $\eta = O(t^2)$. Consequently $B_5(g) = O(t)$, and $F = O(t^6)$ uniformly. Rescaling proves the assertion. $\square$

Exact model and positive order-six correction. On $\mathbb{Z}/5\mathbb{Z}$, choose the Fourier coefficients

$$a_1 = \rho e^{i\pi/10}, \quad a_2 = b e^{i\pi/5}, \quad a_{-1} = \overline{a_1}, \quad a_{-2} = \overline{a_2}.$$

The constant coefficient of $W = \delta + f$ is δ . Put $z = e^{i\pi/10}$, so $z^5 = i$. The six symmetry types of $T(s, t)$ are as follows; multiplicities total 25.

(s, t)	multiplicity	$T(s, t)$
$(0, 0)$	1	$\delta^4 + 2\rho^4 + 2b^4$
$(0, 1)$	4	$2\delta^2 \rho^2 + 2\rho^2 b^2 + b^4$
$(0, 2)$	4	$2\delta^2 b^2 + 2\rho^2 b^2 + \rho^4$
$(1, 1)$	4	$\delta^2 \rho^2 + 2\delta \rho^2 b$
$(2, 2)$	4	$\delta^2 b^2$
$(1, 2)$	8	$\rho^2 b^2 + 2\delta \rho^2 b$

Each entry follows by summing the five bases r in (III.34.4); terms with phase $z^5 = i$ cancel with conjugates. Squaring these entries and summing with their multiplicities proves

$$\|f\|_{U^3}^8 = 8\rho^8 + 16\rho^6 b^2 + 48\rho^4 b^4 + 16\rho^2 b^6 + 8b^8, \quad (\text{III.35.3})$$

$$\begin{aligned} Q_3(\delta + f) &= \delta^8 + 24\delta^4(\rho^4 + b^4) + 16\delta^3 \rho^4 b + 96\delta^2 \rho^4 b^2 + 48\delta^2 \rho^2 b^4 \\ &\quad + 32\delta \rho^4 b^3 + \|f\|_{U^3}^8. \end{aligned} \quad (\text{III.35.4})$$

Choose

$$b = \kappa \delta^{-1} u^2, \quad \kappa = 2^{5/4}/12,$$

and choose the unique positive ρ making (III.35.3) equal u^8 . Existence follows for small u from strict increase in ρ . The expansion is

$$\rho^4 = \frac{u^4}{2\sqrt{2}} - 2^{-3/4} u^2 b^2 + O_\delta(u^8).$$

The order-six coefficient in (III.35.4) is therefore

$$-24 \cdot 2^{-3/4} \kappa^2 + 4\sqrt{2} \kappa = \frac{2^{3/4}}{3}.$$

All other terms are $O_\delta(u^8)$. The real weights are in $[0, 1]$ for small u , since $\rho, b \rightarrow 0$. $\square$

[write] **6 October 2026; sources 21–27.** *The coefficient is now determined* by sources 21 and 23, independently: $2^{3/4}/3$ is optimal, uniformly over all nontrivial groups of exponent five, with an absolute $O(u^8)$ remainder for $u \leq c\delta$ (source 21’s Theorem F.26.1, source 23’s Theorem III.58.1; Section III.49). “The coefficient $2^{3/4}/3$ is a lower bound, not a claim of optimality” stands as this source’s statement. On $\mathbb{Z}/5\mathbb{Z}$ the next coefficient is $20/9$ (source 21’s Theorem III.55.1); the universal eighth-order coefficient is open (item III.27).

The same upper argument applies whenever a largest character has order 5, even if the ambient group, still of order coprime to 6, has characters of other orders. Hence an asymptotically optimal family for the coefficient C_* cannot eventually use an order-five largest pair: its rigidity conclusion has $b \asymp t^{1/3}$ and $\eta \asymp t^{4/3}$, while the order-five penalty would force $\eta \gg b^2 \asymp t^{2/3}$.

III.35.1 The fixed-group classification

For a fixed nontrivial G of order coprime to 6, let $M_{\delta,G}(u)$ denote the weighted supremum with that group held fixed. If the exponent of G is not 5, its dual has a character of order at least 7. For order greater than 8, the positive two-harmonic model in (III.34.17)–(III.34.18) applies without change.

Here is a complete uncentered certificate for order 7, which is the only additional small case needed. Take $a_{\pm 1} = \rho$, $a_{\pm 2} = b$, and $a_0 = \delta$ in $\mathbb{Z}/7\mathbb{Z}$. Under sign changes and interchange of s, t , the following ten types exhaust all 49 slots:

(s, t)	multiplicity	$T(s, t)$
$(0, 0)$	1	$\delta^4 + 2\rho^4 + 2b^4$
$(0, 1)$	4	$2\delta^2\rho^2 + 2\rho^2b^2$
$(0, 2)$	4	$2\delta^2b^2 + \rho^4$
$(0, 3)$	4	$b^4 + 2\rho^2b^2$
$(1, 1)$	4	$\delta^2\rho^2 + 2\delta\rho^2b$
$(1, 2)$	8	$2\delta\rho^2b$
$(1, 3)$	8	ρ^2b^2
$(2, 2)$	4	δ^2b^2
$(2, 3)$	8	0
$(3, 3)$	4	$2\rho b^3$

Squaring and summing gives

$$\begin{aligned} \|f\|_{U^3}^8 &= 8\rho^8 + 48\rho^4b^4 + 32\rho^2b^6 + 8b^8, \\ Q_3(\delta + f) &= \delta^8 + 24\delta^4(\rho^4 + b^4) + 16\delta^3\rho^4b + 96\delta^2\rho^4b^2 + \|f\|_{U^3}^8. \end{aligned}$$

Thus order seven adds only $32\rho^2b^6$ to the centered norm formula and to the final norm term of the cube count.

At $b = \beta_*\delta^{-1/3}u^{4/3}$, choose ρ so that the actual norm equals u . The extra term is $O_\delta(u^{10})$ in the eighth power and changes ρ^4 by $O_\delta(u^6)$. It therefore changes the cube count only by $O_\delta(u^6)$, which is smaller than the secondary scale $u^{16/3}$. All these models lift along a character to the original fixed group, preserving normalized averages.

Consequently, if the exponent of G is not 5,

$$M_{\delta,G}(u) = \delta^8 + 6\sqrt{2}\delta^4u^4 + C_*\delta^{8/3}u^{16/3} + O_\delta(u^{17/3}).$$

If G has exponent 5, the first correction has sharp exponent 6, and its coefficient remains to be determined. The trivial group is excluded because its only centered function is zero.

[write] **6 October 2026; sources 21–27.** The coefficient is now determined: it is $2^{3/4}/3$, uniformly (sources 21 and 23; note after Theorem III.35.2).

III.36 Exact supersaturation and rigidity for polynomial colorings

Source 14, delivered Section 6.

The dimension obstruction for low-degree polynomial colorings has a stronger counting form. Once the dimension exceeds the obstruction budget, the least possible number of symmetrically colored progressions can be determined exactly, at each prescribed center and in total. Equality also has a geometric consequence: the permitted directions at each center form a linear subspace. The classical ingredients are Warning’s second theorem and the minimum-support theorem for polynomial functions on a finite field. We prove the ingredients needed here; their application gives an exact extremal answer in the specified polynomial-coordinate model. Neither the polynomial minimum-support theorem nor the norm construction is claimed as new.

This is a related finite-field problem, rather than a direct replacement for Gowers’s Corollary 8.3. That corollary concerns monochromatic four-term progressions in colorings of an interval. A symmetrically colored four-term progression may have two different colors, in the order a, b, b, a . The repository’s `HasMonochromaticAP` and `corollary_8_3` definitions encode the former condition; the finite-field symmetry condition below needs a separate definition.

III.36.1 The extremal question and its exact answer

Let $\mathbb{F}_q$ have characteristic p , and assume

$$k = 2m \geq 4, \quad p > k, \quad r \geq 1, \quad 1 \leq D_\ell \leq k \quad (1 \leq \ell \leq r).$$

The restriction concerns the characteristic, not just the cardinality q of the field. Put

$$s_\ell = \lceil D_\ell/2 \rceil, \quad B = \sum_{\ell=1}^r s_\ell^2.$$

Consider polynomial maps $Q = (Q_1, \dots, Q_r) : \mathbb{F}_q^d \rightarrow \mathbb{F}_q^r$ with $\deg Q_\ell \leq D_\ell$. At a center z , define

$$S_z(Q) = \left\{ h \in \mathbb{F}_q^d : Q(z + (i - \tfrac{1}{2})h) = Q(z - (i - \tfrac{1}{2})h) \text{ for } 1 \leq i \leq m \right\}. \quad (\text{III.36.1})$$

The zero direction is included. For $h \neq 0$, the resulting k points are distinct because $p > k$. Let

$$A_k^{\text{sym}}(Q) = \#\{(a, h) \in \mathbb{F}_q^d \times (\mathbb{F}_q^d \setminus \{0\}) : Q(a + ih) = Q(a + (k - 1 - i)h) \ (0 \leq i < m)\}.$$

This is a count of ordered pairs (a, h) ; no quotient by changes of orientation or by different parametrizations is implicit.

Theorem III.36.1 (Exact polynomial-coloring supersaturation). *For every admissible Q and every center z ,*

$$|S_z(Q)| \geq q^{\max\{d-B, 0\}}. \quad (\text{III.36.2})$$

There is one admissible Q attaining equality simultaneously at every center. Consequently

$$\min_Q A_k^{\text{sym}}(Q) = q^d (q^{\max\{d-B, 0\}} - 1). \quad (\text{III.36.3})$$

If $d \geq B$ and equality holds in (III.36.2) at a particular z , then $S_z(Q)$ is an $\mathbb{F}_q$ -linear subspace of dimension $d - B$. Equality in (III.36.3) therefore implies this subspace conclusion at every center.

[write] **6 October 2026; sources 21–27.** Source 22 extends the pointwise statement to every equality pattern and every degree at a prescribed centre (its Theorem III.61.5 and Corollary III.62.2, which recovers the low-degree case here), and shows that the quintic and sextic pointwise budget for symmetric four-term progressions is exactly eight (Section III.62).

The previous existence threshold $d \leq B$ is recovered by asking when the right side of (III.36.3) is zero. Above the threshold, the conclusion strengthens the first Chevalley–Warning argument, which supplies only $p - 1$ nonzero directions over a prime field. In dimension $d = B + e$, the exact unavoidable number is $q^e - 1$ at every center.

[write] **6 October 2026; sources 11–14.** *Credit.* The threshold case (the minimum is 0 exactly when $d \leq B$) is source 10’s Theorem III.30.3, the “separate project report” of Section F.19; it is reproved here. New: the exact counts above the threshold over every $\mathbb{F}_q$ with $p > k$, the subspace geometry of equality, and the prime-field gap. Source 10’s quintic question ($4 \leq d \leq 8$) is not answered (item III.19).

III.36.2 A self-contained minimum-support lemma

A polynomial in n variables is called reduced if its degree in each individual variable is at most $q - 1$. Replacing X_i^q by X_i repeatedly gives the unique reduced polynomial representing a given polynomial function on $\mathbb{F}_q^n$. Reduction cannot increase total degree. Uniqueness follows by induction from the fact that a nonzero one-variable polynomial of degree less than q cannot vanish at every element of $\mathbb{F}_q$.

Lemma III.36.2 (Minimum support). *Let F be a nonzero reduced polynomial on $\mathbb{F}_q^n$ with total degree at most $R < n(q - 1)$. Write $R = a(q - 1) + b$, with $0 \leq b < q - 1$. Then*

$$\#\{x : F(x) \neq 0\} \geq (q - b)q^{n-a-1}. \quad (\text{III.36.4})$$

In particular, for $0 \leq B \leq n$, degree at most $B(q - 1)$ and nonempty support imply support size at least q^{n-B} .

Proof. Successively take the leading coefficient in the last variable. More explicitly, if $F(X', X_n) = \sum_{j=0}^t F_j(X')X_n^j$ and $F_t \neq 0$, then above every point where F_t is nonzero, at least $q - t$ values of X_n make F nonzero. Induction on the number of variables therefore produces exponents $e_1, \dots, e_n$, each between 0 and $q - 1$, such that

$$\sum_i e_i \leq R, \quad \#\text{supp } F \geq \prod_i (q - e_i).$$

These exponents form a monomial of F , obtained by the successive leading-coefficient choices.

It remains to minimize the product under the displayed constraints. We may increase some e_i until their sum is R , since this only decreases the product. If $0 < e_i \leq e_j < q - 1$, transferring one unit from e_i to e_j changes the corresponding two factors from $(q - e_i)(q - e_j)$ to $(q - e_i + 1)(q - e_j - 1)$, decreasing their product. Repeating this operation leaves at most one exponent strictly between 0 and $q - 1$. The minimum is attained by a exponents equal to $q - 1$, one equal to b , and the others equal to zero. Its value is exactly $(q - b)q^{n-a-1}$. For $B = n$, the final assertion is simply that nonempty support has at least one point. $\square$

The next equality case is the $b = 0$ part of the classical minimum-weight Reed–Muller classification; see Leducq’s account [57]. A short slicing proof is included to make the geometric conclusion of Theorem III.36.1 independent of a coding-theory black box.

Lemma III.36.3 (Rigidity of minimum support). *If a nonzero reduced polynomial of degree at most $B(q-1)$ on $\mathbb{F}_q^n$ has exactly q^{n-B} nonzero values, where $0 \leq B \leq n$, its support is an affine subspace of dimension $n-B$.*

Proof. The cases $B=0$ and $B=n$ are immediate. Set $e = n - B \geq 1$ and write $S = \text{supp } F$. An invertible affine change of variables, followed by reduction, preserves the degree bound and the support size. Fix any family of q parallel affine hyperplanes and take them to be $X_1 = c$. Let t be the number of these hyperplanes missing S .

If $1 \leq t \leq q-2$, then

$$F(X) = \prod_{c: S \cap \{X_1=c\}=\emptyset} (X_1 - c) G(X)$$

as a polynomial. To see the factorization, the restriction to a missing hyperplane is a reduced polynomial that vanishes on the entire field grid, hence is the zero polynomial; the factor theorem applies to each distinct c . We have

$$\deg G \leq B(q-1) - t = (B-1)(q-1) + (q-1-t).$$

On each of the other $q-t$ hyperplanes, G has nonempty support and its nonzero values occur exactly where F is nonzero. Applying Lemma III.36.2 in $n-1$ variables gives at least $(t+1)q^{e-1}$ points on each such hyperplane. Consequently

$$|S| \geq (q-t)(t+1)q^{e-1} > q^e,$$

because $(q-t)(t+1) - q = t(q-t-1) > 0$. This is impossible. Thus every parallel family has either no missing member or exactly $q-1$ missing members.

In a family with no missing member, applying the minimum-support bound to every restriction shows that each hyperplane contains at least q^{e-1} points of S . The total is q^e , so every member contains exactly q^{e-1} points.

Some family must have a missing member. Otherwise every affine hyperplane would contain exactly q^{e-1} points of S . Choose $x \notin S$ and count pairs (H, y) with H a hyperplane through x and $y \in H \cap S$. There are $(q^n - 1)/(q - 1)$ hyperplanes through x , whereas each $y \neq x$ lies with x in $(q^{n-1} - 1)/(q - 1)$ hyperplanes. The two counts would give

$$q^{e-1}(q^n - 1) = q^e(q^{n-1} - 1),$$

which would imply $q^{e-1} = q^e$. Hence a family has exactly one member meeting S , and S is contained in a hyperplane.

Take that hyperplane to be $X_1 = 0$. Uniqueness of reduced representatives gives

$$F(X_1, X') = (1 - X_1^{q-1})G(X'), \quad G(X') = F(0, X').$$

Because $G \neq 0$, the degree of the right side is $q-1 + \deg G$, so $\deg G \leq (B-1)(q-1)$. Its support still has size q^e . Induction with $n-1$ variables and budget $B-1$ shows that this support is an affine e -dimensional subspace of the hyperplane. $\square$

Corollary III.36.4 (Warning's second theorem, with an equality case). *Suppose polynomials $F_1, \dots, F_t$ on $\mathbb{F}_q^n$ have a nonempty common zero set Z , and their total degree sum is at most $B \leq n$. Then $|Z| \geq q^{n-B}$. If equality holds, Z is an affine subspace of dimension $n-B$.*

Proof. The polynomial function

$$I(X) = \prod_{j=1}^t (1 - F_j(X)^{q-1})$$

is the indicator of Z . Its degree is at most $B(q-1)$, and it is nonzero since Z is nonempty. Reduce I and apply the two lemmas. $\square$

The numerical inequality is precisely the classical second theorem of Warning [40]. Its equality statement here is a consequence of the classical minimum-support classification.

III.36.3 Odd Taylor coefficients and the sharp construction

For fixed z , define the coefficient polynomials by

$$Q_\ell(z + TH) = \sum_{j=0}^{D_\ell} C_{\ell,j,z}(H) T^j.$$

Each nonzero $C_{\ell,j,z}$ is homogeneous of degree j in H . This definition uses polynomial expansion and requires no division by factorials.

Lemma III.36.5 (The reflection equations). *Under the hypotheses above,*

$$S_z(Q) = \{h : C_{\ell,2j-1,z}(h) = 0 \text{ for all } 1 \leq \ell \leq r, 1 \leq j \leq s_\ell\}.$$

Proof. Put $\alpha_i = i - \frac{1}{2}$. The α_i are nonzero and their squares are pairwise distinct: for $i \neq j$ both $i - j$ and $i + j - 1$ are nonzero in characteristic $p > k$. Subtracting the evaluations at $T = \alpha_i$ and $T = -\alpha_i$, then dividing by $2\alpha_i$, gives

$$\sum_{j=1}^{s_\ell} C_{\ell,2j-1,z}(h) (\alpha_i^2)^{j-1} = 0 \quad (1 \leq i \leq m).$$

Since $s_\ell \leq m$, the first s_ℓ equations have an invertible Vandermonde matrix. They are equivalent to vanishing of all the listed coefficients. The converse follows by substitution. $\square$

Proof of Theorem III.36.1. The coefficient polynomials vanish at $h = 0$ and have total degree sum at most

$$\sum_{\ell=1}^r \sum_{j=1}^{s_\ell} (2j-1) = \sum_{\ell=1}^r s_\ell^2 = B.$$

For $d \geq B$, Corollary III.36.4 therefore gives $|S_z(Q)| \geq q^{d-B}$. For $d < B$, the zero direction alone gives the asserted bound. If $d \geq B$ and equality holds, the same corollary says that $S_z(Q)$ is an affine subspace. It contains zero, so it is a linear subspace.

For sharpness, for each positive odd t let $N_t : \mathbb{F}_q^t \rightarrow \mathbb{F}_q$ be the norm from $\mathbb{F}_{q^t}$ to $\mathbb{F}_q$ in a fixed linear coordinate basis. It is a homogeneous polynomial of degree t , and $N_t(u) = 0$ if and only if $u = 0$. These facts follow, for example, by writing

$$N_t(u_1, \dots, u_t) = \prod_{a=0}^{t-1} \left(\sum_{i=1}^t u_i \beta_i^{q^a} \right),$$

where the β_i form a field basis. Frobenius permutes the factors, so the coefficients lie in $\mathbb{F}_q$.

Use disjoint blocks $u_\ell^{(2j-1)} \in \mathbb{F}_q^{2j-1}$, altogether B coordinates, and set

$$Q_\ell(u) = \sum_{j=1}^{s_\ell} N_{2j-1}(u_\ell^{(2j-1)}). \quad (\text{III.36.5})$$

This is the odd norm-layer construction used in the recent coloring literature [33]. Its degree is at most D_ℓ . Suppose a direction has a nonzero block in coordinate ℓ , and let t be its largest nonzero block index. In $Q_\ell(z + Th)$ the coefficient of T^t is exactly $N_t(h_\ell^{(t)}) \neq 0$: larger blocks have zero direction and contribute constants, whereas smaller blocks have degree less than t . Since t is odd, the reflection-coefficient lemma rules out this direction. Thus the map on $\mathbb{F}_q^B$ has $S_z = \{0\}$ at every center.

For $d \geq B$, use this map on the first B coordinates and ignore the remaining $d - B$ coordinates. Its symmetric-direction set is exactly the kernel of that projection, with size q^{d-B} at every center. For $d < B$, restrict the map along a linear injection $\mathbb{F}_q^d \hookrightarrow \mathbb{F}_q^B$; every nonzero direction remains nonzero, so $S_z = \{0\}$.

Finally, the change of variables $z = a + (k - 1)h/2$ is a bijection between start-direction pairs and center-direction pairs. Hence

$$A_k^{\text{sym}}(Q) = \sum_{z \in \mathbb{F}_q^d} (|S_z(Q)| - 1).$$

Summing the pointwise bound and using the construction proves the exact minimum. If the sum attains its lower bound, every nonnegative pointwise excess is zero, proving the final equality assertion. $\square$

There is also an adaptive version: if at a particular center some of the odd coefficient polynomials vanish identically, their degrees should be omitted. With

$$B_z = \sum_{\ell, j: C_{\ell, 2j-1, z} \neq 0} (2j - 1),$$

the same proof gives $|S_z(Q)| \geq q^{\max\{d-B_z, 0\}}$. In particular, equality at the universal bound in a dimension $d \geq B$ forces $B_z = B$.

III.36.4 A sharp gap in the first supercritical dimension

Over a prime field, homogeneity and the elementary Chevalley–Warning divisibility give considerably more than an increment of one above the minimum when $d = B + 1$.

Proposition III.36.6 (First-supercritical gap and stability). *Work over $\mathbb{F}_p$ with $p > k$ and $d = B + 1$. At every center,*

$$|S_z(Q)| = p \quad \text{or} \quad |S_z(Q)| \geq p^2. \quad (\text{III.36.6})$$

Both values p and p^2 occur for admissible maps, simultaneously at all centers in their respective examples.

Let E be the number of centers with $|S_z(Q)| > p$. Then

$$A_k^{\text{sym}}(Q) \geq p^d(p - 1) + (p^2 - p)E. \quad (\text{III.36.7})$$

In particular, if the left side is at most $p^d(p-1) + \varepsilon p^d(p^2-p)$, at least $(1-\varepsilon)p^d$ centers have exactly one projective direction of reflection symmetry: their $S_z(Q)$ is a one-dimensional linear subspace.

Proof. The defining coefficient equations are homogeneous, so their nonzero common zeros form scalar orbits of size $p-1$. Thus $|S_z(Q)| \equiv 1 \pmod{p-1}$. The degree sum is at most $B < d$, so $p \mid |S_z(Q)|$ by Chevalley–Warning. For completeness, this divisibility follows by summing the indicator polynomial from Corollary III.36.4: its total degree is less than $d(p-1)$, and every monomial has a coordinate exponent less than $p-1$. Its sum in that coordinate is zero in $\mathbb{F}_p$, including exponent zero, where the sum is $p=0$.

Writing $|S_z(Q)| = pt$, the two congruences yield $t \equiv 1 \pmod{p-1}$. Since $t \geq 1$, either $t = 1$ or $t \geq p$. This proves (III.36.6). The norm construction controlling B coordinates leaves one free coordinate and gives p directions. Restrict that avoiding norm map to a $(B-1)$ -dimensional subspace, then ignore two coordinates in the $d = B+1$ dimensional domain; this gives exactly p^2 directions. All degrees remain within the same upper bounds. Summing the local gap gives (III.36.7); the last assertion uses the equality geometry in Theorem III.36.1. $\square$

The prime-field restriction in this proposition is deliberate. Warning’s second theorem over $\mathbb{F}_q$ gives the exact minimum q^{d-B} , but the elementary divisibility proof just used supplies divisibility by the characteristic p . It does not by itself justify replacing every occurrence of p in the gap argument by a nonprime q .

III.36.5 Examples and boundaries

For scalar degree at most three or four and $k = 4$, the budget is $B = 4$. Consequently

$$\min_Q A_4^{\text{sym}}(Q) = \begin{cases} 0, & d \leq 4, \\ q^d(q^{d-4} - 1), & d \geq 4. \end{cases}$$

At $q = 5, d = 6$, this is $5^6(5^2 - 1) = 375\,000$ ordered nontrivial progressions. The example is $Q = x_0 + N_3(x_1, x_2, x_3)$ with two additional free coordinates. Each center has exactly 25 symmetric directions, including zero. At $q = 5, d = 5$, every center has either exactly five or at least twenty-five symmetric directions.

The equality subspaces may vary with the center. For example, $Q(x, y) = x + y^2$ has degree at most two, $B = 1$, and

$$S_{(a,b)}(Q) = \{(u, v) : u + 2bv = 0\}.$$

There are q directions at every center, attaining the minimum in dimension two; the one-dimensional kernel changes with b . Consequently equality does not, in general, imply that Q factors through a single fixed linear projection of dimension B .

The higher-degree question remains distinct. For a scalar quintic and $k = 4$, write $P(z + TH) = \sum_{j=0}^5 C_{j,z}(H)T^j$. Reflection at the two parameters $1/2$ and $3/2$ is equivalent to

$$C_{1,z} = \frac{9}{16}C_{5,z}, \quad C_{3,z} = -\frac{5}{2}C_{5,z}.$$

Thus the odd coefficients need not vanish individually, and the argument giving the low-degree budget $B = 9$ cannot be used as an avoidance construction for four-term progressions. Conversely, the same equations are equivalent to

$$C_{3,z} + \frac{40}{9}C_{1,z} = 0, \quad C_{5,z} - \frac{16}{9}C_{1,z} = 0.$$

These two polynomials have degrees at most three and five, respectively. Their degree sum is therefore at most $3 + 5 = 8$, giving the necessary bound $d \leq 8$ for avoidance. The cubic construction gives $d \geq 4$. Nothing in the present supersaturation theorem determines the maximum within $4 \leq d \leq 8$.

Proposition III.36.7 (Failure of the first norm extension in characteristic five). *For every prime p , the degree- p field norm $N : \mathbb{F}_{p^p} \rightarrow \mathbb{F}_p$ is constant on a nontrivial affine $\mathbb{F}_p$ -line. In particular, at $p = 5$ a quintic norm in five variables does not avoid symmetrically colored four-term progressions. Adding independent lower-degree norm blocks does not remove this counterexample.*

Proof. Choose a root θ of $X^p - X - 1$ in an algebraic closure. Then $\theta^p = \theta + 1$, and induction gives $\theta^{p^j} = \theta + j$ for $0 \leq j \leq p$. The first p elements are distinct, and the p th iterate returns to θ . Thus $X^p - X - 1$ is irreducible over $\mathbb{F}_p$, and θ lies in $\mathbb{F}_{p^p}$. The points $\theta + a$, for $a \in \mathbb{F}_p$, are precisely its Frobenius conjugates, and

$$N(\theta + a) = \prod_{b \in \mathbb{F}_p} (\theta + a + b) = \theta^p - \theta = 1.$$

They form the required affine line. In a direct sum of norm blocks, fix every other block and vary only this line in the degree- p block; the total color remains constant. $\square$

This obstruction concerns a proposed construction, not every polynomial of degree five. It therefore does not improve the upper bound eight or disprove the possibility of a different five-dimensional coloring.

III.37 Quantitative examples and what the refinements measure

Source 14, delivered Section 7.

III.37.1 Convergence of the sharp cube model

The sharp secondary term comes from a competition between two effects. A small coefficient at twice the dominant frequency creates a positive quintic contribution, linear in the secondary amplitude. Keeping the centered U^3 norm fixed requires decreasing the dominant amplitude, producing a quartic penalty in that same secondary amplitude. In normalized variables the competition is exactly

$$4\sqrt{2}tb - 48b^4.$$

The balance $b \asymp t^{1/3}$ explains the fractional power: the outer quartic scale t^4 is multiplied by $t^{4/3}$. The coefficient is determined by the penalty as well as the positive correlation. Retaining only the coordinate axes of the Fourier identity misses the eight extra penalty terms and therefore misses the sharp coefficient.

For the positive two-harmonic model, define

$$\mathcal{R}_\delta(u) = \frac{Q_3(\delta + f_u) - \delta^8 - 6\sqrt{2}\delta^4 u^4}{\delta^{8/3} u^{16/3}}.$$

At $\delta = 1/4$, the exact polynomial formulas give the following values, rounded from 90-digit decimal arithmetic.

u	$\mathcal{R}_{1/4}(u)$	$C_* = (9/4)^{1/3}$
10^{-2}	1.3532657043	1.3103706971
10^{-3}	1.3123554554	1.3103706971
10^{-4}	1.3104628080	1.3103706971
10^{-6}	1.3103708955	1.3103706971
10^{-9}	1.3103706971	1.3103706971




Figure III.37.1: The exact two-harmonic model at density $1/4$. The left panel displays the normalized secondary coefficient; the right displays its positive excess over C_* on logarithmic axes. These are evaluations of the proved polynomial identities, not fitted evidence for the theorem. The model's error after normalization is $O_\delta(u^{4/3})$; the uniform upper theorem allows a larger remainder because it controls all functions.

On an exponent-five group, the additional penalty is quadratic in b , rather than quartic at the relevant scale. The balance there is $b \asymp t$, giving a correction t^6 . This explains why a torsion assumption is substantive even after order three has been excluded. It also distinguishes a proved exponent from an optimal coefficient: the order-six exponent is sharp, while its best coefficient remains a concrete open problem.

III.37.2 A sharp stability example and a computable moment gain

Take $A = \mathbb{F}_{101} \setminus \{0\}$. Then $m = 100$ and

$$E(A) = 990\,100 = (1 - 0.0099)m^3.$$

The nearest coset is the whole field, at distance one. Thus

$$\frac{|A \Delta \mathbb{F}_{101}|}{|A|} = 0.01 = \tau(0.0099).$$

Taking a Cartesian product with full additional coordinates preserves both normalized ratios. This gives a transparent sharpness example in every cube dimension.

The same example shows why a denominator depending only on the sizes of the cross-sections can already improve a moment bound. A section of size 100 in a prime-coordinate space cannot be a coordinate coset, since such a coset has size 101^a . The cardinality defect in Corollary II.44.3 is 0.01, yielding an energy deficit 0.0099 without inspecting individual points. If every nonempty section has this same normalized defect, the eighth-arrangement lower bound gains the factor

$$(1 - 0.0099)^{-6} > 1.06$$

over the cubic-profile denominator. The gain is local and explicit; whether it survives subsequent steps depends on the parameter regime of the full argument.

III.37.3 How fast polynomial supersaturation grows

For one cubic or quartic color coordinate and symmetric four-term progressions, $B = 4$. Over $\mathbb{F}_5$ the theorem gives the exact values below. Directions include zero in the middle column and exclude it in the final column.

Dimension d	Minimum directions at each center	Minimum ordered nontrivial four-APs
4	1	0
5	5	12 500
6	25	375 000
7	125	9 687 500

The norm construction attains every row. Adding one dimension beyond the threshold does not merely force one progression somewhere: it forces four nonzero directions at every center. In the first supercritical dimension the next possible local count is at least twenty-five including zero, so a near-minimum global count strongly limits the number of exceptional centers. The equality kernels can still vary with the center, as the quadratic example in Section III.36 demonstrates.

III.38 A finite certificate for the centered cube expansion

Source 14, delivered Appendix A.

Associate to a cube vertex $\omega \in \{0, 1\}^3$ the row $(1, \omega_1, \omega_2, \omega_3)$. A vertex in a chosen support is called a *coloop* if deleting its row lowers the rank of the rows. If a support has a coloop, its product of centered functions has zero average: one may average the corresponding linear form independently of all the others.

Here is the full support classification. Ranks are over the rationals; the relevant nonzero four-by-four minors are ± 1 and ± 2 , so the same independence argument is valid over every finite abelian group of odd order.

Size	Total	With a coloop	Remaining	Contribution
0	1	0	1	Constant term
1	8	8	0	Zero
2	28	28	0	Zero
3	56	56	0	Zero
4	70	58	12	Parallelogram moment S
5	56	48	8	Five-point moment B_5
6	28	0	28	At most u^6 each
7	8	0	8	At most u^7 each
8	1	0	1	u^8

For four vertices, the dependent supports are the six faces and six diagonal parallelograms. Every other four-vertex support has full rank. For five vertices, the complement is a three-vertex set. Up to cube symmetries, its three pairwise Hamming distances are $(1, 1, 2)$, $(1, 2, 3)$, or $(2, 2, 2)$. These types have respectively 24, 24, and 8 realizations. In the first two types the five-vertex support has a coloop; the final type is the set of three neighbors of a cube vertex. Its complement is a cube-symmetric copy of $\{000, 100, 010, 001, 111\}$.

For example, the complements of $\{000, 100, 010\}$ and $\{000, 100, 111\}$ have coloops 110 and 011, respectively. In the remaining case the five forms satisfy the unique full-support dependence

$$2L_{000} - L_{100} - L_{010} - L_{001} + L_{111} = 0, \quad L_\omega = x + \omega \cdot h.$$

Orthogonality then gives the Fourier expression for B_5 used in Equation (III.34.6). Generalized cube Cauchy–Schwarz, obtained by successive Cauchy–Schwarz in the cube directions, bounds a support of size j by $\|f\|_{U_3}^j$ after inserting constant ones at the other vertices. This proves the exact expansion and the two remainder bounds in Equation (III.34.7).

The file `verify_certificates.py` enumerates all 255 nonempty supports using exact rational row reduction, checks the coloop counts and the minor values, and verifies the five-vertex description. It also checks all 25 and 49 Fourier slots in the order-five and order-seven tables with exact integer polynomial coefficients. Thus no floating point computation is needed for these certificates.

III.39 Exact fourth- and fifth-order cube corrections in every dimension

Source 16, delivered Section 8.

Gowers’s Lemma 3.10 bounds the number of cubes by applying the triangle inequality to a constant plus a balanced function. The resulting binomial bound includes many terms which vanish identically. We determine all terms through degree five in every dimension. The answer includes a three-torsion contribution which first occurs in dimension four.

III.39.1 Notation and the expansion

Let G be a finite abelian group, with normalized averages, and let $f : G \rightarrow \mathbb{R}$ satisfy $\mathbb{E}f = 0$. Put $m = 2^d$ and define

$$\mathcal{Q}_d(g) = \mathbb{E}_{x, h_1, \dots, h_d \in G} \prod_{v \in \{0,1\}^d} g(x + v \cdot h).$$

For real g , this is $\|g\|_{U^d}^m$. Write $a_\chi = \widehat{f}(\chi)$, using the normalization $f(x) = \sum_{\chi \in \widehat{G}} a_\chi \chi(x)$ and additive notation for $\widehat{G}$. Thus $a_{-\chi} = \overline{a_\chi}$ and $a_0 = 0$. Set

$$Q(f) = \sum_{\chi \in \widehat{G}} |a_\chi|^4 = \|f\|_{U^2}^4, \quad (\text{III.39.1})$$

$$\begin{aligned} Q_2(f) &= \sum_{2\chi=0} |a_\chi|^4, \\ D(f) &= \sum_{\chi \in \widehat{G}} a_{2\chi} a_\chi a_{-\chi}^3, \\ D_3(f) &= \sum_{3\chi=0} a_\chi^4 a_{-\chi}. \end{aligned} \quad (\text{III.39.2})$$

All four quantities are real. The first two are nonnegative; the last two can have either sign. In each sum the trivial character contributes zero. Define

$$\begin{aligned} P_d &= \frac{6^d - 2 \cdot 4^d + 2^d}{8}, & T_d &= \frac{8^d - 3 \cdot 6^d + 3 \cdot 4^d - 2^d}{24}, \\ F_d &= 4T_d, & H_d &= \frac{10^d - 4 \cdot 8^d + 6 \cdot 6^d - 4 \cdot 4^d + 2^d}{24}. \end{aligned} \quad (\text{III.39.3})$$

Theorem III.39.1 (The low-order cube expansion). *Let $d \geq 2$, and let $f : G \rightarrow \mathbb{R}$ have mean zero. The coefficients of t, t^2, t^3 in $\mathcal{Q}_d(\delta + tf)$ vanish. Over every finite abelian group, the coefficient of t^4 is*

$$\delta^{m-4} (P_d Q(f) + T_d Q_2(f)). \quad (\text{III.39.4})$$

If G has odd order and $d \geq 3$, then

$$\begin{aligned} \mathcal{Q}_d(\delta + tf) &= \delta^m + P_d \delta^{m-4} Q(f) t^4 \\ &\quad + \delta^{m-5} (F_d D(f) + H_d D_3(f)) t^5 + R_d(\delta, t, f). \end{aligned} \quad (\text{III.39.5})$$

Writing $\tau = \|f\|_{U^d}$, the remainder satisfies

$$|R_d(\delta, t, f)| \leq \sum_{k=6}^m \binom{m}{k} |\delta|^{m-k} |t|^k \tau^k. \quad (\text{III.39.6})$$

For $d = 2$ the exact formula is $\mathcal{Q}_2(\delta + tf) = \delta^4 + t^4 Q(f)$.

[write] **6 October 2026; sources 15–20.** *Credit.* An independent third derivation of source 09's Theorem III.21.1 (on every finite abelian group; this F_d, H_d are source 09's Q_d, S_d , this P_d is source 03's R_d), also proved by source 18 (Theorem III.40.1); the quartic part is source 03's Theorem III.2.6 and source 08's. The restriction of the quintic formula to odd order is not needed: the intake found the same formula on $\mathbb{Z}/2$, $\mathbb{Z}/4$, $\mathbb{Z}/6$, $\mathbb{Z}/8$ and $(\mathbb{Z}/2)^2$, as source 09's all-groups theorem implies. The $\mathbb{Z}/3\mathbb{Z}$ example (200 and 352) is also in source 18's Section V.77. *Notation:* D here is source 04's B_5 and source 09's M_5 ; D_3 is source 09's $T_{3,5}$.

d	P_d	T_d	F_d	H_d
2	1	0	0	0
3	12	2	8	0
4	100	40	160	16
5	720	520	2080	480
6	4816	5600	22400	8960

For example, in dimension four and odd order,

$$\mathcal{Q}_4(\delta + tf) = \delta^{16} + 100\delta^{12}Q(f)t^4 + \delta^{11}(160D(f) + 16D_3(f))t^5 + R_4. \quad (\text{III.39.7})$$

The term $16D_3(f)$ cannot generally be dropped: it is present on groups with three-torsion. On groups of order coprime to six, only Q and D occur through degree five.

Example III.39.2 (A nonzero three-torsion correction). On $\mathbb{Z}/3\mathbb{Z}$, let $f = (2, -1, -1)$. Its two non-trivial normalized Fourier coefficients are both one, so $Q(f) = D(f) = D_3(f) = 2$. Formula (III.39.7) begins

$$\mathcal{Q}_4(\delta + tf) = \delta^{16} + 200\delta^{12}t^4 + 352\delta^{11}t^5 + O(t^6).$$

Omitting the order-three term would instead give 320 as the quintic coefficient. This is an exact algebraic distinction, not a limiting effect.

III.39.2 Vertex systems and the quartic term

For a vertex subset $S = \{v_1, \dots, v_k\} \subseteq \{0, 1\}^d$, put

$$M_f(S) = \mathbb{E}_{x,h} \prod_{i=1}^k f(x + v_i \cdot h).$$

Fourier orthogonality gives

$$M_f(S) = \sum_{\substack{\xi_1 + \dots + \xi_k = 0 \\ \sum_i v_{i,j} \xi_i = 0 \ (1 \leq j \leq d)}} \prod_{i=1}^k a_{\xi_i}. \quad (\text{III.39.8})$$

The constraints are the kernel equations of the integer matrix with columns $(1, v_i)$. Consequently the Taylor coefficient of degree k is $\delta^{m-k} \sum_{|S|=k} M_f(S)$.

Lemma III.39.3. *Any one, two, or three distinct Boolean vertices have a maximal affine minor equal to 1 or -1 . Thus $M_f(S) = 0$ for $1 \leq |S| \leq 3$.*

Proof. Reflect coordinates to put one vertex at zero. Coordinate reflection preserves the average: reflect coordinate j by replacing x with $x + h_j$ and h_j with $-h_j$. For three vertices, the two remaining nonzero distinct binary columns have at least two distinct nonzero row patterns from 10, 01, 11. Any two such patterns have determinant ± 1 . The cases of one or two vertices are immediate. The corresponding linear evaluation map onto $G^{|S|}$ is therefore surjective, so its outputs are independent uniform variables. Their product average is $(\mathbb{E}f)^{|S|} = 0$. $\square$

Lemma III.39.4 (Four-vertex classification). *For four distinct Boolean vertices, exactly one of the following applies:*

1. *They form a real parallelogram, and $M_f(S) = Q(f)$.*
2. *Their affine matrix has full rank and an affine minor ± 1 , and $M_f(S) = 0$.*
3. *They form an affine $\mathbb{F}_2$ plane which is not a real parallelogram, and $M_f(S) = Q_2(f)$.*

The numbers of configurations in the first and third cases are P_d and T_d , respectively.

Proof. Reflect one vertex to zero, leaving three binary columns. Their rational rank is two or three. In rank two, project injectively onto two suitable coordinates. The four images are all of the two-cube; hence the original vertices form a parallelogram with primitive affine relation $(1, 1, -1, -1)$. Lemma III.39.3 supplies a unimodular subsystem, so the frequency kernel is exactly $(\chi, \chi, -\chi, -\chi)$, in a suitable ordering. Formula (III.39.8) gives $Q(f)$.

In rank three, each three-by-three binary determinant is $0, \pm 1$, or ± 2 . A minor ± 1 makes the evaluation map surjective and gives zero. The sole determinant-two binary pattern, up to row and column permutations, has rows 110, 101, 011. Adding any other nonzero binary row produces a determinant-one minor with two of these rows. Therefore, if no minor ± 1 exists, every coordinate row is 000, 110, 101, or 011, and all three nonzero types occur. The kernel equations force all four frequencies to be the same χ with $2\chi = 0$. Since a_χ is real on such characters, their contribution is $Q_2(f)$.

To count parallelograms, choose their common midpoint. If r coordinates vary between the endpoints of an opposite pair, there are 2^{d-r} choices for the fixed coordinates and 2^{r-1} distinct opposite pairs about that midpoint. Thus their number is

$$\sum_{r=2}^d \binom{d}{r} 2^{d-r} \binom{2^{r-1}}{2} = \frac{6^d - 2 \cdot 4^d + 2^d}{8} = P_d.$$

There are $2^d(2^d - 1)(2^d - 2)/24$ affine two-planes in $\mathbb{F}_2^d$. Subtracting P_d gives T_d . □

III.39.3 Five-vertex circuits and the three-torsion term

Lemma III.39.5 (Five-vertex classification in odd order). *Suppose that G has odd order. For a five-element vertex subset S , the only possibly nonzero contributions are the following:*

1. *A rational affine circuit with relation $2v_0 + v_4 = v_1 + v_2 + v_3$, contributing $D(f)$. There are F_d such subsets.*
2. *An affinely independent configuration which is a signed combinatorial embedding of $\{0000, 1110, 1101, 1011, 0111\}$, contributing $D_3(f)$. There are H_d such subsets.*

Every other five-element subset has $M_f(S) = 0$.

Proof. A rational affine r -plane contains at most 2^r Boolean vertices: projection onto r suitable coordinates is injective on that plane. Hence the affine dimension of S is three or four.

First suppose it is three, and project injectively onto three coordinates. Every five vertices of a three-cube contain a unimodular four-vertex subsystem. Indeed, their affine rank modulo two is three because an affine $\mathbb{F}_2$ plane has only four points. Some affine minor is therefore odd; all such minors have absolute value at most two, so this minor is ± 1 . The frequency kernel is the free rank-one kernel of the primitive rational relation, with no additional torsion. Extra ambient coordinate rows impose no further constraints, because the same integer relation annihilates each of them.

If four vertices form a parallelogram, the unique relation has a zero coefficient at the remaining vertex. Its frequency is consequently zero, and $a_0 = 0$ makes the contribution vanish. Otherwise the relation has full support. The three-cube has precisely eight five-point circuits, namely a vertex, its three neighbors, and its antipode. To see exhaustion, classify the omitted triple by

its Hamming distances: the only types are $(1, 1, 2)$, $(1, 2, 3)$, and $(2, 2, 2)$. For the first two types take the representatives $\{000, 100, 010\}$ and $\{000, 100, 111\}$; their complements contain the parallelograms $\{001, 011, 101, 111\}$ and $\{001, 010, 101, 110\}$, respectively. The third type leaves one of the stated circuits. Its primitive relation is $2v_0 + v_4 = v_1 + v_2 + v_3$. Thus the frequency kernel is $(2\chi, -\chi, -\chi, -\chi, \chi)$, and (III.39.8) gives $D(f)$.

For the count, v_0 is the unique vertex with coefficient two and v_4 is also determined; v_1, v_2, v_3 have $3!$ possible orderings. In any coordinate the relation allows eight binary row patterns: two constants and three complementary pairs of nonconstant patterns. All three nonconstant types must occur to give affine dimension three. By inclusion–exclusion, the number of subsets is

$$\frac{2^d}{6}(4^d - 3 \cdot 3^d + 3 \cdot 2^d - 1) = F_d.$$

Now suppose the affine dimension is four. Some five-by-five affine minor is nonzero and has absolute value at most three. In fact, replacing the four binary coordinate rows by sign rows multiplies the determinant by 2^4 , while Hadamard’s inequality bounds the resulting determinant by $5^{5/2}$; hence the original integer determinant has absolute value at most $\lfloor 5^{5/2}/16 \rfloor = 3$.

A minor of absolute value one or two gives a surjective evaluation map on an odd-order group and therefore contributes zero. Otherwise every nonzero affine minor has absolute value three. Modulo three, a row basis can be extended from the nonzero constant row; rank five would therefore give an affine minor nonzero modulo three. The augmented rank modulo three is exactly four: any four of these affinely independent vertices remain independent modulo three, since their nonzero affine minors have absolute value at most two. Its one-dimensional mod-three kernel has full support. Its five nonzero coefficients are ± 1 , and their sum is zero modulo three. Four coefficients therefore have one sign and the remaining coefficient has the other sign.

Let v_0 be the uniquely distinguished vertex. Reflect it to zero. Every coordinate row of the other four vertices has either zero or three ones. All four weight-three patterns must occur, or the rational affine rank would be smaller than four. This is precisely the configuration in the statement. Its frequency equations force $\xi_i = \chi$ for $1 \leq i \leq 4$, $\xi_0 = -\chi$, and $3\chi = 0$. Thus its contribution is $D_3(f)$. In particular, only three-torsion occurs, not characters of order nine: the equations, or the adjugate of a determinant-three subsystem, give $3\chi = 0$.

There are now ten permitted coordinate patterns: two constants and four complementary pairs of nonconstant patterns. All four nonconstant types must occur. The distinguished vertex is unique and the other four have $4!$ possible orderings. Their number is therefore

$$\frac{2^d}{24}(5^d - 4 \cdot 4^d + 6 \cdot 3^d - 4 \cdot 2^d + 1) = H_d.$$

This exhausts both possible rational affine dimensions. □

Proof of Theorem III.39.1. Expand the product defining $\mathcal{Q}_d(\delta + tf)$ by its vertex subsets. Lemmas III.39.3–III.39.5 determine every coefficient through degree five. For a subset S of cardinality k , apply Gowers–Cauchy–Schwarz with f on S and the constant function one on its complement. This gives $|M_f(S)| \leq \tau^k$. Summing the remaining terms gives (III.39.6). □

Remark III.39.6 (The embedding count). If $\{ \begin{smallmatrix} n \\ r \end{smallmatrix} \}$ denotes a Stirling number of the second kind, the number of signed combinatorial r -subcubes in a d -cube is

$$E_{d,r} = 2^{d-r} \left\{ \begin{matrix} d+1 \\ r+1 \end{matrix} \right\}.$$

Partition the varying coordinates into r nonempty blocks, copy or complement one latent variable on each block, and fix the remaining coordinates. The extra element in the Stirling number marks the block of fixed coordinates. The formulas above become $P_d = E_{d,2}$, $T_d = 2E_{d,3}$, $F_d = 8E_{d,3}$, and $H_d = 16E_{d,4}$.

III.39.4 Consequences, sharp moment bounds, and limits

Corollary III.39.7 (A refined version of the cube-counting upper bound). *Let $d \geq 3$, let G have odd order, let $A \subseteq G$ have density δ , and put $f = 1_A - \delta$. If $\|f\|_{U^d}^m \leq \alpha$, then*

$$\begin{aligned} \mathcal{Q}_d(1_A) &\leq \delta^m + P_d \delta^{m-4} Q(f) + \delta^{m-5} (F_d D(f) + H_d D_3(f)) \\ &\quad + \sum_{k=6}^m \binom{m}{k} \delta^{m-k} \alpha^{k/m}. \end{aligned}$$

If only α is available, monotonicity gives $Q(f) \leq \alpha^{4/m}$. The following estimates bound the fifth-order terms as well. In particular, at fixed d and positive δ , the leading error is of order $\alpha^{4/m}$ instead of the order $\alpha^{1/m}$ appearing in the direct Minkowski bound. The exact expression retains more useful information whenever Q , D , or D_3 can be estimated separately.

Proposition III.39.8 (Sharp bounds for the quintic moment). *For real mean-zero functions on odd-order groups,*

$$|D(f)| \leq 2^{-1/4} Q(f)^{5/4}, \quad |D_3(f)| \leq 2^{-1/4} Q(f)^{5/4}.$$

If the group has no three-torsion, then

$$|D(f)| \leq 2^{-1/2} Q(f)^{5/4}.$$

The constants are sharp in their stated classes.

Proof. Pair nonzero frequencies into $\{\chi, -\chi\}$. Put $b_i = |a_\chi|$ for one representative of each pair. Doubling permutes these pairs, say by σ , because the group has odd order. Then

$$|D(f)| \leq 2 \sum_i b_i^4 b_{\sigma(i)}, \quad Q(f) = 2 \sum_i b_i^4.$$

Bounding each $b_{\sigma(i)}$ by $(\sum_j b_j^4)^{1/4}$ proves the first inequality. The same argument restricted to three-torsion proves the bound for D_3 . Equal positive Fourier coefficients on an order-three character pair attain both constants.

If three-torsion is absent, σ has no fixed points. Normalize $\sum_i b_i^4 = 1$ and let $u = \max_i b_i$. If $u^4 \leq 1/2$, the sum is at most $u \leq 2^{-1/4}$. Otherwise put $v = (1 - u^4)^{1/4}$. The largest coordinate is sent to another coordinate, at most v , while all other coordinates are sent to values at most u . Hence

$$\sum_i b_i^4 b_{\sigma(i)} \leq u^4 v + v^4 u \leq 2^{-1/4}.$$

For the last inequality, the function $(r^4 + r)/(r^4 + 1)^{5/4}$ has its maximum $2^{-1/4}$ at $r = 1$: the sign of its derivative is the sign of

$$(1 - r)(r^6 + r^5 + r^4 + 5r^3 + r^2 + r + 1).$$

Restoring the normalization proves the second constant. Equal positive Fourier coefficients on the four nontrivial characters of an order-five quotient attain it. $\square$

[write] **6 October 2026; sources 15–20.** *Credit.* The bound with $2^{-1/4}$ is source 04’s bound on B_5 (Section III.3); the constant $2^{-1/2}$ without three-torsion, sharp, is new in this report (the intake’s adversarial search found no exceedance).

Remark III.39.9 (Higher torsion and the scope of the improvement). The low-order proof suggests organizing subsequent coefficients by the integer kernels and Smith normal forms of affine Boolean vertex matrices. A full-rank k -vertex subsystem has a nonzero affine minor of absolute value at most $\lfloor k^{k/2}/2^{k-1} \rfloor$. Primes larger than this cannot produce a new torsion kernel for that subsystem. For $k = 4, 5$ this bound gives exactly the possible primes two and three.

An explicit family in dimension $r \geq 3$ consists of the $r + 1$ vertices $0, \mathbf{1} - e_1, \dots, \mathbf{1} - e_r$. Its affine determinant is $r - 1$ and its correlation is

$$\sum_{(r-1)\chi=0} a_{-\chi} a_{\chi}^r.$$

The coordinate equations force all nonzero-vertex frequencies to equal χ and then force $(r - 1)\chi = 0$. This family need not exhaust the configurations or locate the first occurrence of a given prime.

The results above are explicit local improvements to the cube-counting step. They do not, by themselves, improve a global quantitative Szemerédi bound. Their proofs are independent of the accompanying finite calculations; the exact computations serve only as reproducible checks of classification counts, signs, and normalization.

III.39.5 Optimality of the exponent for sets

The fourth-order term also gives an obstruction for genuine subsets of cyclic groups, rather than only for weighted functions.

Proposition III.39.10 (Sharp order of the small-uniformity error). *Fix $d \geq 2$ and put $m = 2^d$. There are odd cyclic groups G_n and subsets $A_n \subseteq G_n$ of density exactly $\delta = 1/3$ such that, writing $\alpha_n = \|1_{A_n} - \delta\|_{U^d}^m$, one has*

$$\alpha_n \longrightarrow 0, \quad \frac{\mathcal{Q}_d(1_{A_n}) - \delta^m}{\alpha_n^{4/m}} \longrightarrow c_d > 0.$$

Consequently a bound with a uniform error $o(\alpha^{4/m})$ is impossible, even at this fixed density and on odd cyclic groups.

Proof. On $G_0 = \mathbb{Z}/3\mathbb{Z}$, take $f_0 = (2, -1, -1)$ and $w_t = 1/3 + tf_0$, where $0 < t < 1/3$. This is a function into $[0, 1]$ of mean $1/3$. The two nontrivial Fourier coefficients of f_0 are one, so $Q(f_0) = 2$. Let $B_d = \|f_0\|_{U^d}^m > 0$. Theorem III.39.1 gives

$$\begin{aligned} \mathcal{Q}_d(w_t) &= \delta^m + 2P_d\delta^{m-4}t^4 + O_d(t^5), \\ \mathcal{Q}_d(w_t - \delta) &= B_d t^m. \end{aligned} \tag{III.39.9}$$

For $d = 2$ the first identity is exact without the error term.

We explain carefully how to realize these two moments by sets. Let $p \neq 3$ be an odd prime and independently select each $(x, y) \in G_0 \times \mathbb{Z}/p\mathbb{Z}$ with probability $w_t(x)$. For a random cube, two distinct vertices have the same second coordinate with probability $1/p$, because their difference is a

nonzero binary linear form with a coefficient 1 or -1 . Away from these collisions, the selected-point variables at its vertices are independent. Thus, for $j = 1, 2$, the expected value of

$$T_1 = \mathcal{Q}_d(1_A), \quad T_2 = \mathcal{Q}_d(1_A - \delta)$$

differs from the corresponding weighted moment in (III.39.9) by at most $2\binom{m}{2}/p$.

Likewise, products associated with two cubes are independent unless they share a selected-point variable. A union bound over their m^2 pairs of vertices gives

$$\text{Var}(T_j) \leq \frac{2m^2}{p}.$$

The density has expectation δ and variance at most $1/(12p)$. Consequently, for any $\varepsilon > 0$ and sufficiently large p , a realization exists whose density differs from δ by at most ε and whose two moments differ from their weighted targets by at most 2ε . This follows directly from Chebyshev's inequality and a union bound.

Change at most an ε fraction of the points to make the cardinality exactly p , hence the density exactly $1/3$. A telescoping product and uniformity of each vertex marginal show that this changes each T_j by at most $m\varepsilon$. Choose $\varepsilon = t^{m+1}$ and then choose $p = p(t)$ sufficiently large. The resulting set, still denoted by A_t , satisfies

$$\begin{aligned} \mathcal{Q}_d(1_{A_t}) - \delta^m &= 2P_d\delta^{m-4}t^4 + O_d(t^5), \\ \|1_{A_t} - \delta\|_{U^d}^m &= B_d t^m + O_d(t^{m+1}). \end{aligned}$$

Taking any sequence $t \downarrow 0$ and using $G_0 \times \mathbb{Z}/p\mathbb{Z} \cong \mathbb{Z}/(3p)\mathbb{Z}$ proves the claim, with

$$c_d = 2P_d\delta^{m-4}B_d^{-4/m} > 0.$$

□

[write] **6 October 2026; sources 15–20.** *Credit.* Sharpness of the exponent $4/2^d$ for sets is source 03's Corollary III.2.11, source 04's Proposition III.3.10, source 09's Proposition III.22.1 with Corollary III.22.5, and source 18's Theorem III.40.12; new here is only the packaging for sets of exact density $1/3$ in odd cyclic groups. At $d = 2$ the ratio is identically one, so that case is trivial.

III.40 Exact fourth and fifth order cube counts

Source 18, delivered Section 3.

Gowers' Lemma 3.10 applies the triangle inequality to the cube norm. The resulting estimate is convenient, but it treats the constant and balanced parts as if the first mixed term could already be nonzero. Centering actually eliminates all terms of degrees one, two, and three. We now compute the next two degrees exactly, in every dimension and on every finite abelian group. The answer distinguishes ordinary affine circuits from the extra relations created by two- and three-torsion.

Throughout this section, G is a finite abelian group, averages are normalized, and the dual group $\hat{G}$ is written additively. Our Fourier convention is

$$\hat{f}(\gamma) = \mathbb{E}_{x \in G} f(x) \overline{\gamma(x)}, \quad \hat{G}[p] = \{\gamma \in \hat{G} : p\gamma = 0\}.$$

For a real function g , define

$$\mathcal{C}_d(g) = \mathbb{E}_{x, h_1, \dots, h_d \in G} \prod_{v \in \{0,1\}^d} g(x + v \cdot h) = \|g\|_{U^d}^{2^d}.$$

Thus $|G|^{d+1} \mathcal{C}_d(1_A)$ counts parameterized cubes in A , including cubes with repeated vertices. Real-valuedness is essential to the simple identical-function expansion used below; complex cube norms carry the usual alternating conjugations.

III.40.1 The coefficient theorem

For $d \geq 3$, put $m = 2^d$ and define

$$R_d = \frac{6^d - 2 \cdot 4^d + 2^d}{8}, \quad B_d = \frac{2^d(2^d - 1)(2^d - 2)}{24}, \quad (\text{III.40.1})$$

$$Q_d = \frac{8^d - 3 \cdot 6^d + 3 \cdot 4^d - 2^d}{6}, \quad S_d = \frac{10^d - 4 \cdot 8^d + 6 \cdot 6^d - 4 \cdot 4^d + 2^d}{24}. \quad (\text{III.40.2})$$

For a real mean-zero function $f : G \rightarrow \mathbb{R}$, set

$$V_4(f) = \sum_{\gamma \in \widehat{G}} |\widehat{f}(\gamma)|^4 = \|f\|_{U^2}^4, \quad (\text{III.40.3})$$

$$E_2(f) = \sum_{\gamma \in \widehat{G}[2]} \widehat{f}(\gamma)^4, \quad (\text{III.40.4})$$

$$T_5(f) = \sum_{\gamma \in \widehat{G}} \widehat{f}(2\gamma) \widehat{f}(\gamma) \widehat{f}(-\gamma)^3, \quad (\text{III.40.5})$$

$$E_3(f) = \sum_{\gamma \in \widehat{G}[3]} \widehat{f}(-\gamma) \widehat{f}(\gamma)^4. \quad (\text{III.40.6})$$

These quantities are real: replacing γ by $-\gamma$ conjugates each summand. Moreover, $V_4(f)$ and $E_2(f)$ are nonnegative. In the latter sum, every Fourier coefficient is real because $2\gamma = 0$. The zero character contributes zero to all four expressions.

Theorem III.40.1 (Exact expansion through degree five). *Let G be any finite abelian group, $d \geq 3$, $m = 2^d$, and let $f : G \rightarrow \mathbb{R}$ have mean zero. For all real δ, t ,*

$$\begin{aligned} \mathcal{C}_d(\delta + tf) &= \delta^m + \delta^{m-4} t^4 (R_d V_4(f) + (B_d - R_d) E_2(f)) \\ &\quad + \delta^{m-5} t^5 (Q_d T_5(f) + S_d E_3(f)) + \mathcal{R}_d(\delta, t; f), \end{aligned} \quad (\text{III.40.7})$$

where, writing $\eta = \|f\|_{U^d}$,

$$|\mathcal{R}_d(\delta, t; f)| \leq \sum_{j=6}^m \binom{m}{j} |\delta|^{m-j} |t|^j \eta^j. \quad (\text{III.40.8})$$

In particular, the coefficients of t, t^2, t^3 vanish identically.

[write] **6 October 2026; sources 15–20.** *Credit.* The coefficients are those of source 09's Theorem III.21.1 on every finite abelian group ($Q_d = 2^d S(d+1, 4)$, $S_d = 2^d S(d+1, 5)$ with Stirling numbers, as source 09 writes them), also derived independently by source 16 (Theorem III.39.1, odd order for the quintic part); the quartic part is source 03's Theorem III.2.6 and source 08's. So this is the fourth independent derivation in

the report; source 18 could see none of the other three (sources 09 and 16 were not in the repository at its pin). New in method: the Smith-form and Hadamard-minor argument showing that only the primes 2 and 3 enter through order five (Corollary III.40.4). *Notation:* R_d is source 03's R_d (source 04's, 09's and 16's P_d), $B_d - R_d$ is T_d ; this Q_d is a five-circuit count (= source 09's Q_d = source 16's F_d), not a cube count.

Remark III.40.2 (Small torsion and low dimensions). If G has no nonzero element killed by 2, then $E_2(f) = 0$. If G has no nonzero element killed by 3, then $E_3(f) = 0$. The fifth order identity therefore needs no exclusion of two-torsion. For $d = 2$ the whole expansion reduces to the familiar exact identity

$$\mathcal{C}_2(\delta + tf) = \delta^4 + t^4 \|f\|_{U^2}^4.$$

The first coefficient values are

d	R_d	B_d	Q_d	S_d
3	12	14	8	0
4	100	140	160	16
5	720	1240	2080	480

The twelve four-point circuits and eight five-point circuits in the ordinary three-dimensional cube are classical; they are listed, for example, by Beerenwinkel, Pachter, and Sturmfels [74]. Here the aim is an explicit all-dimensional cube-moment formula, including the contributions of small torsion and a norm-controlled remainder. The circuit interpretation itself is not a novelty claim.

III.40.2 Fourier constraints and integer minors

For a set $W = \{v_1, \dots, v_s\} \subseteq \{0, 1\}^d$, let

$$M_W = \begin{pmatrix} 1 & \cdots & 1 \\ v_1 & \cdots & v_s \end{pmatrix}.$$

Fourier inversion and character orthogonality give

$$\mathbb{E}_{x,h} \prod_{v \in W} f(x + v \cdot h) = \sum_{\substack{\gamma_1, \dots, \gamma_s \in \widehat{G} \\ M_W \gamma = 0}} \prod_{i=1}^s \widehat{f}(\gamma_i). \quad (\text{III.40.9})$$

The matrix acts by integer multiplication in the abelian group $\widehat{G}$. Rational rank alone does not determine this kernel: the integer invariant factors must also be controlled.

Lemma III.40.3 (Small minors of Boolean affine matrices). *For a Boolean affine matrix M_W , any $k \times k$ minor containing the row of ones has absolute value at most $1, 1, 1, 2, 3$ for $k = 1, 2, 3, 4, 5$, respectively. Every three distinct columns have a 3×3 minor equal to ± 1 . Every five distinct columns have a 4×4 minor equal to ± 1 .*

Proof. For a minor containing the row of ones, reflect each selected coordinate in which the first selected vertex equals one. At the matrix level this replaces a coordinate row by the row of ones minus that row. The first selected vertex now has all selected coordinates zero. Subtracting its column from the other columns reduces the determinant, up to sign, to that of a binary $(k-1) \times (k-1)$ matrix.

If D is a binary $r \times r$ matrix, the matrix

$$H = \begin{pmatrix} 1 & \mathbf{1}^\top \\ \mathbf{1} & J - 2D \end{pmatrix}$$

has entries in $\{-1, 1\}$ and determinant $(-2)^r \det D$. Hadamard's inequality gives

$$|\det D| \leq \frac{(r+1)^{(r+1)/2}}{2^r}.$$

For $r = 1, 2, 3, 4$, taking integer parts gives $1, 1, 2, 3$.

Three distinct points of $\mathbb{F}_2^d$ have affine rank three, because an affine line contains only two points. Extend the nonzero row of ones to a row basis over $\mathbb{F}_2$ and select three independent columns. The resulting minor is odd and has absolute value at most one, so it is ± 1 .

Similarly, five distinct points have affine rank at least four over $\mathbb{F}_2$, because an affine plane contains only four points. Extending the row of ones to four independent rows and selecting four independent columns produces an odd minor of absolute value at most two. It too must equal ± 1 . $\square$

We use the elementary Smith normal form principle that unimodular integer row and column operations induce automorphisms over every abelian group. If a full-column-rank integer matrix has invariant factors $1, \dots, 1, p$, its kernel over an abelian group H is a copy of $H[p]$. For prime p , its coordinates are the multiples of a vector spanning the one-dimensional kernel modulo p .

Corollary III.40.4 (Only two small primes can occur). *For four distinct cube vertices, the rational rank is three or four. In rank three the nonzero Smith invariants are $1, 1, 1$; in rank four they are $1, 1, 1, D$ with $D \in \{1, 2\}$.*

For five distinct cube vertices, the rational rank is four or five. In rank four the nonzero invariants are $1, 1, 1, 1$; in rank five they are $1, 1, 1, 1, D$ with $D \in \{1, 2, 3\}$.

Proof. The unit minors in Lemma III.40.3 make the third determinantal divisor equal to one for a four-point set and the fourth determinantal divisor equal to one for a five-point set. In full column rank, extend the row of ones to a rational row basis. A nonzero maximal minor containing that row has absolute value at most two or three. The final determinantal divisor divides this minor. The stated lists follow. $\square$

III.40.3 The four-point coefficient

Lemma III.40.5 (Four-point classification). *For four distinct Boolean cube vertices, exactly the following possibilities can contribute to (III.40.9):*

1. *An ordinary rectangle, with primitive relation $u + v = w + z$, contributes $V_4(f)$.*
2. *Four rationally independent vertices whose XOR is zero contribute $E_2(f)$.*

All other four-point sets have zero correlation. There are R_d sets of the first type and $B_d - R_d$ of the second.

Proof. Suppose first that the rational rank is three. Its unique affine dependence has nonzero coefficients, since three distinct vertices are independent. The positive and negative coefficients split as two and two: a cube vertex cannot be a strict convex combination of other cube vertices. Write

$$au + bv = cw + ez, \quad a + b = c + e,$$

with all coefficients positive. Reflect coordinates so that $u = 0$. In every nonconstant coordinate, $v = 1$ and a nonempty proper subset of $\{c, e\}$ sums to b . Affine dimension two requires both possible nonconstant patterns. Therefore $b = c = e$ and then $a = b$. The primitive dependence is $u + v = w + z$. The unit Smith invariants show that the Fourier kernel consists exactly of the tuples $(\gamma, \gamma, -\gamma, -\gamma)$, giving $V_4(f)$.

In rational rank four, the kernel is trivial unless the final Smith invariant is two. Modulo two, the unique dependence must involve all four columns because every three are independent. It is therefore the XOR relation. The Fourier kernel is exactly

$$\{(\gamma, \gamma, \gamma, \gamma) : 2\gamma = 0\},$$

giving $E_2(f)$. A trivial Fourier kernel gives a product containing $\hat{f}(0) = 0$.

To count rectangles, there are six binary coordinate patterns satisfying $u + v = w + z$. The two degenerate matchings each give 4^d ordered tuples, and their intersection gives 2^d . Every nondegenerate rectangle has eight labelings of this form: choose which opposite pair is on the left and order the two pairs. This gives R_d in (III.40.1).

An unordered four-set has XOR zero precisely when it is an affine plane in $\mathbb{F}_2^d$. Choose its origin, then an ordered basis, and divide by the four origins and six ordered bases of one plane. The count is

$$\frac{2^d(2^d - 1)(2^d - 2)}{4 \cdot 6} = B_d.$$

Of these, R_d are ordinary rectangles, leaving $B_d - R_d$ of the second type. □

III.40.4 The five-point coefficient

Lemma III.40.6 (Ordinary five-point circuits). *Every five-point circuit of the real affine Boolean cube has, after relabeling and common sign, the primitive relation*

$$2u + v = x + y + z. \tag{III.40.10}$$

There are exactly Q_d such circuits, and each contributes $T_5(f)$ over every finite abelian group. Every rational rank-four five-point set that is not a circuit contributes zero.

Proof. In rational rank four, Corollary III.40.4 gives a unit rank-four minor. The character kernel over any $\hat{G}$ is therefore generated by the primitive integer dependence vector, without additional torsion parameters. If this vector has a zero coefficient, its support consists of four points forming a rectangle. The remaining Fourier coordinate is zero, so the correlation vanishes.

Otherwise, all five coefficients are nonzero. Their signs split as two and three, again because every cube vertex is extreme. Write the positive dependence as

$$au + bv = cx + ey + qz, \quad a + b = c + e + q,$$

where all five coefficients are positive. Reflect coordinates so that $u = 0$. Every nonconstant coordinate has $v = 1$ and corresponds to a nonempty proper subset of $\{c, e, q\}$ whose sum is b . Affine dimension three requires at least three distinct such patterns.

For three positive numbers, three distinct proper nonempty subsets can have a common sum only if all three numbers are equal. Indeed, equal-sum subsets form an antichain. A three-element antichain of proper nonempty subsets of a three-element set is either the three singletons or the three pairs: a singleton and a pair omitting it leave no third subset incomparable with both. Equality of the three singleton sums or three pair sums forces $c = e = q$. Thus $b = c$ or $b = 2c$, and $a = 2c$ or $a = c$, respectively. Interchanging u, v if necessary proves (III.40.10).

Its Fourier kernel is

$$(\gamma_u, \gamma_v, \gamma_x, \gamma_y, \gamma_z) = (2\gamma, \gamma, -\gamma, -\gamma, -\gamma),$$

which gives (III.40.5).

For the count, the coefficient-two vertex and the same-sign coefficient-one vertex are uniquely distinguished; the other three vertices have $3!$ orderings. Each coordinate row satisfying (III.40.10) is one of eight patterns: two constants and three nonconstant types together with their complements. After reflecting to $u = 0$, the three types are

$$(0, 1, 1, 0, 0), \quad (0, 1, 0, 1, 0), \quad (0, 1, 0, 0, 1).$$

Together with the row of ones, they are independent. Rational rank four holds exactly when all three types occur. Inclusion-exclusion therefore gives

$$Q_d = \frac{8^d - 3 \cdot 6^d + 3 \cdot 4^d - 2^d}{3!}.$$

□

Lemma III.40.7 (The additional three-torsion circuits). *Among rationally independent five-point subsets, exactly S_d can give nonzero centered correlation. Each has a unique dependence modulo three with coefficient pattern $(-1, 1, 1, 1, 1)$, up to common sign and relabeling, and contributes $E_3(f)$. The subsets with a nontrivial invariant factor two contribute zero.*

Proof. Let D be the last Smith invariant in Corollary III.40.4. If $D = 1$, the Fourier kernel is trivial and the correlation is zero.

If $D = 2$, the unique dependence modulo two has even support because the first row consists of ones. Its support has size at least four and at most five, so it has size four. The Smith kernel consists of two-torsion multiples of this vector. One Fourier coordinate is consequently zero, and mean zero again kills the correlation.

Suppose $D = 3$. Any four columns are rationally independent, and a nonzero four-by-four minor containing the row of ones has absolute value one or two. They are therefore independent modulo three. The unique dependence modulo three involves all five columns, with coefficients in $\{1, -1\}$. Their sum is zero modulo three. After a common sign, exactly one is negative and the other four are positive. The Fourier kernel is parameterized by

$$(-\gamma, \gamma, \gamma, \gamma, \gamma), \quad 3\gamma = 0,$$

and its correlation is $E_3(f)$.

For counting, label the distinguished negative vertex u and the others $v_1, \dots, v_4$. Coordinatewise the relation is

$$-u + v_1 + v_2 + v_3 + v_4 = 0 \pmod{3}.$$

There are ten binary row patterns: two constants and four nonconstant types together with their complements. After reflection to $u = 0$, the nonconstant rows are $(0, \mathbf{1} - e_j)$ for $1 \leq j \leq 4$. The four vectors $\mathbf{1} - e_j$ are independent over the rationals because $\det(J - I) = -3$. Thus rational rank five holds exactly when all four nonconstant types occur. Inclusion-exclusion gives

$$10^d - 4 \cdot 8^d + 6 \cdot 6^d - 4 \cdot 4^d + 2^d$$

ordered tuples. The distinguished vertex is unique, and the other four have $4!$ orderings. Division by $4!$ proves the formula for S_d . $\square$

Proof of Theorem III.40.1. Expand the m -fold product defining $\mathcal{C}_d(\delta + tf)$. A subset of j vertices contributes $\delta^{m-j}t^j$ times its f -correlation. For $j = 1, 2, 3$, the corresponding affine matrices have unit full-column-rank minors; character orthogonality forces zero Fourier coordinates, so the correlations vanish. Lemmas III.40.5, III.40.6, and III.40.7 give the fourth and fifth coefficients.

For every subset of j vertices, Gowers–Cauchy–Schwarz, applied to f at those vertices and the constant function one at the others, bounds the absolute correlation by $\|f\|_{U^d}^j$. Summing this bound over all subsets with $j \geq 6$ gives (III.40.8). $\square$

Remark III.40.8 (Characteristic two and characteristic three). If G has exponent two, then $T_5(f) = 0$ because $\hat{f}(2\gamma) = \hat{f}(0) = 0$, and $E_3(f) = 0$ as well. Every fifth order term vanishes. If G has exponent three, then $T_5(f) = E_3(f)$, after replacing γ by $-\gamma$ in one sum. The fifth coefficient is then

$$(Q_d + S_d)E_3(f), \quad Q_d + S_d = \frac{10^d - 6 \cdot 6^d + 8 \cdot 4^d - 3 \cdot 2^d}{24}.$$

These identities explain why an odd-order hypothesis suffices for the ordinary fourth coefficient but does not remove the extra fifth order contribution.

III.40.5 A quantitative strengthening of Lemma 3.10

Corollary III.40.9 (Centered cube bound). *Let $g : G \rightarrow \mathbb{R}$ have mean $\delta \geq 0$, set $f = g - \delta$, and suppose $\|f\|_{U^d} \leq \varepsilon$, with $d \geq 3$ and $m = 2^d$. Then*

$$\begin{aligned} \mathcal{C}_d(g) &\leq \delta^m + B_d \delta^{m-4} \varepsilon^4 + (Q_d + S_d) \delta^{m-5} \varepsilon^5 \\ &\quad + \sum_{j=6}^m \binom{m}{j} \delta^{m-j} \varepsilon^j. \end{aligned} \tag{III.40.11}$$

If G has no two-torsion, B_d may be replaced by R_d . If G has no three-torsion, $Q_d + S_d$ may be replaced by Q_d . The stronger signed spectral formula (III.40.7) remains available in every case.

Proof. Use Theorem III.40.1 with $t = 1$. The fourth coefficient is a sum of B_d correlations, each of absolute value at most ε^4 . The nonzero fifth types number $Q_d + S_d$, each of absolute value at most ε^5 . These estimates follow from the same Gowers–Cauchy–Schwarz inequality as the remainder bound. The stated torsion exclusions remove the corresponding terms from the exact formula. $\square$

For a set A of density δ , the balanced function is $f = 1_A - \delta$. In Gowers' normalization, degree- $(d-1)$ α -uniformity means $\|f\|_{U_d}^m \leq \alpha$. Thus (III.40.11) improves the first possible error from the order $\alpha^{1/m}$ supplied by $(\delta + \alpha^{1/m})^m$ to order $\alpha^{4/m}$, at fixed positive density and fixed d . This is a local improvement of the cube-counting lemma. An improvement of the terminal Szemerédi bound would require propagating it through an argument that uses this estimate quantitatively.

III.40.6 Sharpness, including genuine indicator sets

The unrestricted finite-group leading constant B_d is optimal. The first step is an exact weighted example; a Bernoulli approximation then transfers its leading behavior to sets.

Proposition III.40.10 (Exact order-two character example). *Let $\chi : G \rightarrow \{-1, 1\}$ be a nontrivial character, $d \geq 2$, and $m = 2^d$. Then*

$$\mathcal{C}_d(\delta + t\chi) = \frac{(\delta + t)^m + (\delta - t)^m + (2m - 2)(\delta^2 - t^2)^{m/2}}{2m}. \quad (\text{III.40.12})$$

Moreover, $\|t\chi\|_{U^d} = |t|$. At fixed $\delta > 0$,

$$\mathcal{C}_d(\delta + t\chi) = \delta^m + B_d \delta^{m-4} t^4 + O_{d,\delta}(t^6). \quad (\text{III.40.13})$$

Proof. The values of $\chi(x), \chi(h_1), \dots, \chi(h_d)$ are independent uniform signs. Consequently the signs on the cube are the evaluations of a uniformly random affine Boolean function on $\mathbb{F}_2^d$. Of its $2^{d+1} = 2m$ possible affine functions, two are constant and the other $2m - 2$ are balanced. A constant function contributes $(\delta + t)^m$ or $(\delta - t)^m$; a balanced one contributes $(\delta + t)^{m/2}(\delta - t)^{m/2}$. This proves (III.40.12).

Every d -fold cube derivative of a character is one for $d \geq 2$, so its U^d norm is one. In the even polynomial (III.40.12), the quadratic coefficient is zero and the quartic coefficient is

$$\frac{2\binom{m}{4} + (2m - 2)\binom{m/2}{2}}{2m} = \frac{m(m-1)(m-2)}{24} = B_d.$$

This proves the expansion. □

Lemma III.40.11 (Approximation of weights by sets in cube norm). *Fix $d \geq 2$, put $m = 2^d$, and let $G = \mathbb{F}_2^n$, of order N . For every function $w : G \rightarrow [0, 1]$, there is a set $A \subseteq G$ such that*

$$\|1_A - w\|_{U^d} \leq \left(\frac{m-1}{N} \right)^{1/m}. \quad (\text{III.40.14})$$

Proof. Include each $y \in G$ in A independently with probability $w(y)$, and write $z = 1_A - w$. For a cube whose m sites are distinct, independence and centering give

$$\mathbb{E}_A \prod_{v \in \{0,1\}^d} z(x + v \cdot h) = 0.$$

Every other product has absolute value at most one. A collision occurs only if $a \cdot h = 0$ for some nonzero $a \in \mathbb{F}_2^d$. For fixed nonzero a , this has probability $1/N$. A union bound over $m-1$ choices gives

$$0 \leq \mathbb{E}_A \|z\|_{U^d}^m \leq \frac{m-1}{N}.$$

Some realization is at most this average, proving the claim. □

Theorem III.40.12 (Optimal unrestricted leading coefficient for sets). *Fix $d \geq 3$ and $\delta \in (0, 1)$. There are finite groups $G_k = \mathbb{F}_2^{n_k}$ and sets $A_k \subseteq G_k$, with densities $\delta_k \rightarrow \delta$ and balanced norms $\varepsilon_k = \|1_{A_k} - \delta_k\|_{U^d} \rightarrow 0$, such that*

$$\lim_{k \rightarrow \infty} \frac{\mathcal{C}_d(1_{A_k}) - \delta_k^m}{\delta_k^{m-4} \varepsilon_k^4} = B_d. \quad (\text{III.40.15})$$

Therefore B_d in the unrestricted finite-group leading term of (III.40.11) cannot be decreased, even when the functions being counted are indicators.

Proof. Let $t_k > 0$ tend to zero, with $t_k \leq \min(\delta, 1 - \delta)$. Choose $N_k = 2^{n_k}$ sufficiently large that

$$\left(\frac{m-1}{N_k} \right)^{1/m} \leq t_k^5.$$

On G_k , choose a nontrivial sign character χ_k and set $w_k = \delta + t_k \chi_k$. This weight lies in $[0, 1]$ and has mean δ . Lemma III.40.11 gives a set A_k with $z_k = 1_{A_k} - w_k$ satisfying $\|z_k\|_{U^d} \leq t_k^5$.

The mean is bounded by the cube norm, so

$$\delta_k = \delta + \mathbb{E} z_k = \delta + O(t_k^5).$$

The balanced function of A_k is

$$1_{A_k} - \delta_k = t_k \chi_k + z_k - \mathbb{E} z_k.$$

The triangle inequality and its reverse give $\varepsilon_k = t_k + O(t_k^5)$. Also, $\|1_{A_k}\|_{U^d}, \|w_k\|_{U^d} \leq 1$, because both functions take values in $[0, 1]$. Hence

$$|\mathcal{C}_d(1_{A_k}) - \mathcal{C}_d(w_k)| \leq m \|1_{A_k} - w_k\|_{U^d} \leq m t_k^5.$$

Proposition III.40.10 now yields

$$\mathcal{C}_d(1_{A_k}) - \delta_k^m = B_d \delta^{m-4} t_k^4 + O_{d,\delta}(t_k^5).$$

Since $\delta_k^{m-4} \varepsilon_k^4 = \delta^{m-4} t_k^4 (1 + o(1))$, division proves (III.40.15). $\square$

[write] **6 October 2026; sources 15–20.** *Credit.* Sharpness of B_d for sets is source 03’s Corollary III.2.11 (with exact rational densities in its Theorem III.2.12); proved again here by rounding. The best coefficient on odd-order groups, which Section V.78 asks for (“especially in dimension four”), is settled at $d = 3$ by source 04’s Theorem III.3.4 ($c_3^{\text{odd}} = 2^{-1/2}$, note there); open for $d \geq 4$ (item III.1).

This sharpness statement uses groups with two-torsion. It leaves open the best leading coefficient when multiplication by two is invertible, especially for $d \geq 4$. The exact coefficient in that case is $R_d \|f\|_{U^2}^4$, so the remaining optimization concerns the strongest possible comparison between the centered real U^2 and U^d norms.

III.41 Source 20’s chapter: quadratic counting errors at Section 3 [write]

Added in the fourth write, 6 October 2026.

The next seven sections are source 20, a manuscript on the passage from uniformity to counting in Gowers’s Section 3 (Theorem 3.2, Corollary 3.3). It keeps two functions through the last Cauchy–Schwarz step, which gives an error quadratic in the balanced uniformity norm. *Logically the chapter stands beside source 06’s Section III.5*, which proves the quadratic error first, and source 10’s Section III.28 on torsion; source 20 could see neither (only source 04’s manuscript, whose telescope is linear). It is printed here, at the end of Part III, because no number printed earlier may move. Its introduction is Section F.25; its formalization notes, questions, conclusion and source audit are in Part V (Sections V.83–V.86). Its notation is in the table at the head of this Part.

What is proved, in one paragraph. For surjective homomorphisms of finite abelian groups, two marked functions bounded by directional seminorms and one L^2 factor (Theorem III.43.3; source 06’s Theorem III.5.1 for slope systems); the directional-to-global comparison with the optimal constant for every subgroup tuple (Theorem III.44.1) and a collapsed-pair obstruction; quadratic counting with unequal densities (Theorem III.45.2) and, for one set, $|\Lambda_k(1_A) - \delta^k| \leq B_k(\delta)u^2$ (Corollary III.45.4) with B_k strictly below source 06’s coefficient (proved at the write, note there); the sharp pair correction $\sqrt{m-1}$ (Proposition III.46.1); supersaturation (Corollary III.47.1); and sharpness of the exponent 2 for three-term counts of indicator sets of density near $1/2$ on $\mathbb{F}_{29}^{2m}$ (Theorem III.48.3; new in this report). Not claimed: sharpness for $k \geq 4$ or for prime cyclic groups, a substitute for the asymmetric Corollary 3.3, a global bound.

Relation to the formal statements (checked at HEAD 0b4890a1c). `theorem_3_2` and `corollary_3_3` (Sections01_03.lean, lines 258, 276) are exact companions for prime N and $k \leq N$. Nothing in this chapter is formalized.

III.42 Normalized cubes and subgroup directions

Source 20, delivered Section 2.

All groups in this article are finite abelian groups, written additively. Every expectation uses the uniform probability measure on its indicated finite set. We put $\Delta_h f(x) = f(x)f(x+h)$. Complex conjugation is denoted by $\mathcal{C}$, so $\mathcal{C}^0 z = z$ and $\mathcal{C}^1 z = \bar{z}$. The L^p norms also use normalized measure.

Definition III.42.1 (Directional cube seminorm). For subgroups $H_1, \dots, H_s \leq G$, $s \geq 1$, define

$$D_{H_1, \dots, H_s}(f)^{2^s} = \mathbb{E}_{x \in G} \mathbb{E}_{h_r \in H_r} \prod_{\omega \in \{0,1\}^s} \mathcal{C}^{|\omega|} f(x + \omega \cdot h). \quad (\text{III.42.1})$$

Its nonnegative 2^s -th root is intended. When every $H_r = G$, this is $\|f\|_{U^s}$. In particular $\|f\|_{U^1} = |\mathbb{E}f|$.

The direction list is symmetric: a permutation of its entries simply renames the cube coordinates. Directions may coincide or be trivial. For $s = 1$,

$$D_H(f)^2 = \mathbb{E}_{q \in G/H} |\mathbb{E}_{x \in q} f(x)|^2. \quad (\text{III.42.2})$$

Thus a trivial direction is meaningful and should not be discarded.

Lemma III.42.2 (Positivity, mixed Cauchy–Schwarz, and phase invariance). *The quantity in (III.42.1) is nonnegative. For functions $f_\omega : G \rightarrow \mathbb{C}$,*

$$\left| \mathbb{E}_{x \in G} \mathbb{E}_{h_r \in H_r} \prod_{\omega} \mathcal{C}^{|\omega|} f_\omega(x + \omega \cdot h) \right| \leq \prod_{\omega} D_{H_1, \dots, H_s}(f_\omega). \quad (\text{III.42.3})$$

The resulting D satisfies the triangle inequality. If $s \geq 2$ and χ is a character of G , then $D(f\chi) = D(f)$, and in particular $\|f\chi\|_{U^s} = \|f\|_{U^s}$.

Proof. Averaging the last derivative first gives

$$D_{H_1, \dots, H_s}(f)^{2^s} = \mathbb{E}_{h_1, \dots, h_{s-1}} \mathbb{E}_{q \in G/H_s} |\mathbb{E}_{x \in q} \Delta_{h_1} \cdots \Delta_{h_{s-1}} f(x)|^2. \quad (\text{III.42.4})$$

This proves positivity directly, including $s = 1$.

For the mixed inequality introduce independent $t_r^0, t_r^1 \in H_r$ and an independent $z \in G$. Set $h_r = t_r^1 - t_r^0$ and $x = z + \sum_r t_r^0$. The resulting (x, h) is uniformly distributed, so the mixed expression becomes

$$\mathbb{E}_z \mathbb{E}_{t_r^0, t_r^1} \prod_{\omega} \mathcal{C}^{|\omega|} f_{\omega} \left(z + \sum_r t_r^{\omega_r} \right).$$

For fixed z this is a product-space box form. The box Cauchy–Schwarz inequality follows by applying ordinary Cauchy–Schwarz successively to the s coordinates. At coordinate r one regards all other variables as parameters, squares the average over t_r , and introduces its independent copy. After all s steps, the factor belonging to each ω is its diagonal box form to the power $1/2^s$. This argument also proves positivity of each diagonal box form. Applying Hölder in z , with 2^s equal exponents, gives (III.42.3), because the average of each diagonal box form is exactly (III.42.1).

Expanding the diagonal cube for $f + g$ and applying (III.42.3) to its 2^{2^s} terms gives $D(f + g)^{2^s} \leq (D(f) + D(g))^{2^s}$. Finally, an alternating cube of a character is identically one when $s \geq 2$: both its constant exponent and each linear increment exponent cancel. This proves phase invariance. $\square$

For reference, normalized Fourier coefficients are

$$\widehat{f}(\chi) = \mathbb{E}_x f(x) \overline{\chi(x)}.$$

Orthogonality gives Parseval and the fourth-moment identity

$$\sum_{\chi} |\widehat{f}(\chi)|^2 = \|f\|_2^2, \quad \|f\|_{U^2}^4 = \sum_{\chi} |\widehat{f}(\chi)|^4. \quad (\text{III.42.5})$$

The second identity follows by expanding all four factors of the parallelogram and averaging the three independent variables. Their character constraints leave one freely chosen frequency. These normalizations are used throughout the sharpness argument.

III.43 The two-function directional estimate

Source 20, delivered Section 3.

Let Ω be a finite abelian parameter group, and let $L_1, \dots, L_k : \Omega \rightarrow G$ be surjective homomorphisms. Write

$$\Lambda_L(f_1, \dots, f_k) = \mathbb{E}_{z \in \Omega} \prod_{i=1}^k f_i(L_i z), \quad K_j = \ker L_j, \quad H_{ij} = L_i(K_j) \quad (i \neq j). \quad (\text{III.43.1})$$

For a fixed position i , set

$$D_i(f) = D_{(H_{ij})_{j \neq i}}(f). \quad (\text{III.43.2})$$

This is an order $k - 1$ directional seminorm. There is no pairwise independence assumption in this definition.

III.43.1 The last two coordinates have a common quotient

Lemma III.43.1 (Joint fibers). *For distinct i, j , there is an isomorphism*

$$\theta_{ij} : G/H_{ij} \longrightarrow G/H_{ji}$$

such that the uniform law of $(L_i z, L_j z)$ can be sampled as follows: choose $q \in G/H_{ij}$ uniformly, then choose independently a uniform $u \in q$ and a uniform $v \in \theta_{ij}(q)$. Consequently

$$\mathbb{E}_z F(L_i z) Q(L_j z) = \mathbb{E}_q (\mathbb{E}_{u \in q} F(u)) \left(\mathbb{E}_{v \in \theta_{ij}(q)} Q(v) \right). \quad (\text{III.43.3})$$

Proof. Let $J = (L_i, L_j)(\Omega) \leq G \times G$. Both projections of J are surjective. Its horizontal kernel is $H_{ij} \times \{0\}$ and its vertical kernel is $\{0\} \times H_{ji}$. If $(u, v) \in J$, the class of v modulo H_{ji} depends only on the class of u modulo H_{ij} : two choices differ by a horizontal or vertical kernel element. This constructs an isomorphism θ_{ij} ; symmetry supplies its inverse. For each q , the elements of J above q are precisely the Cartesian product $q \times \theta_{ij}(q)$. These products all have the same size. A homomorphism from Ω to J has equal-size fibers, so uniform z induces uniform measure on J . Formula (III.43.3) follows. $\square$

Lemma III.43.2 (One kernel-averaging step). *Let $L : \Omega \rightarrow G$ be surjective, $K = \ker L$, and let $a : G \rightarrow \mathbb{C}$. For any $F : \Omega \rightarrow \mathbb{C}$,*

$$|\mathbb{E}_z a(Lz) F(z)|^2 \leq \|a\|_2^2 \mathbb{E}_{z \in \Omega} \mathbb{E}_{h \in K} F(z) \overline{F(z+h)}. \quad (\text{III.43.4})$$

The expression on the right after $\|a\|_2^2$ is real and nonnegative.

Proof. Disintegrate uniform measure on Ω over the fibers of L . The left-hand average is $\mathbb{E}_{v \in G} a(v) \mathbb{E}_{z \in L^{-1}(v)} F(z)$. Cauchy–Schwarz bounds its square by $\|a\|_2^2$ times the average of the squared fiber means. Two independent points in the same fiber have the form $(z, z+h)$ with h uniform in K , proving the identity and nonnegativity. $\square$

Theorem III.43.3 (Two marked functions, with an L^2 gain). *Assume $k \geq 3$. For any three distinct positions i, j, ℓ and any functions $f_1, \dots, f_k : G \rightarrow \mathbb{C}$,*

$$|\Lambda_L(f_1, \dots, f_k)| \leq D_i(f_i) D_j(f_j) \|f_\ell\|_2 \prod_{r \notin \{i, j, \ell\}} \|f_r\|_\infty. \quad (\text{III.43.5})$$

For $k = 2$, the corresponding assertion is $|\Lambda_L(f_1, f_2)| \leq D_1(f_1) D_2(f_2)$. The numerical constant one is best possible.

Proof. If an unmarked factor is identically zero, there is nothing to prove. By homogeneity, normalize all factors outside $\{i, j, \ell\}$ to have supremum norm at most one. Eliminate the $r = k - 2$ unmarked positions in an order $a_1 = \ell, a_2, \dots, a_r$.

After t eliminations let $h_b \in K_{a_b}$, $1 \leq b \leq t$, and put

$$F_q^{(t)}(z; h) = \Delta_{L_q h_1} \cdots \Delta_{L_q h_t} f_q(L_q z), \quad A_t = \mathbb{E}_{h_1, \dots, h_t, z} \prod_{q \notin \{a_1, \dots, a_t\}} F_q^{(t)}(z; h).$$

Thus $A_0 = \Lambda_L(f)$. Lemma III.43.2, first with L_ℓ , gives $|A_0|^2 \leq \|f_\ell\|_2^2 A_1$. At a later step, the factor $F_{a_{t+1}}^{(t)}$ is invariant under $z \mapsto z + h_{t+1}$ for $h_{t+1} \in K_{a_{t+1}}$, and is one-bounded. Apply the same

averaging and Cauchy–Schwarz argument on the product probability space of z and all previous h_b . It gives $|A_t|^2 \leq A_{t+1}$. Each A_{t+1} so obtained is nonnegative. Iteration yields

$$|\Lambda_L(f)|^{2^r} \leq \|f_\ell\|_2^{2^r} A_r. \quad (\text{III.43.6})$$

Only i, j remain in A_r .

For a fixed list $h = (h_1, \dots, h_r)$ apply Lemma III.43.1 to their derivative functions. Define

$$M_i(h, q) = \mathbb{E}_{x \in q} \Delta_{L_i h_1} \cdots \Delta_{L_i h_r} f_i(x),$$

and define M_j using the corresponding coset $\theta_{ij}(q)$. Then $A_r = \mathbb{E}_{h, q} M_i(h, q) M_j(h, \theta_{ij}(q))$. Cauchy–Schwarz on this joint probability space gives

$$|A_r| \leq (\mathbb{E}_{h, q} |M_i(h, q)|^2)^{1/2} (\mathbb{E}_{h, q} |M_j(h, \theta_{ij}(q))|^2)^{1/2}.$$

The first squared moment is exactly

$$\mathbb{E}_{h_1, \dots, h_r} \mathbb{E}_{x \in G} \mathbb{E}_{t \in H_{ij}} \Delta_{L_i h_1} \cdots \Delta_{L_i h_r} f_i(x) \overline{\Delta_{L_i h_1} \cdots \Delta_{L_i h_r} f_i(x+t)} = D_i(f_i)^{2^{r+1}}.$$

Indeed $L_i h_b$ is uniform in H_{i, a_b} , and these images are independent for different b . The same computation gives $D_j(f_j)^{2^{r+1}}$ for the other moment. Combining with (III.43.6) and taking the 2^r -th root proves (III.43.5). Restoring the normalized supremum norms gives the stated homogeneous form.

For $k = 2$, apply Cauchy–Schwarz directly to (III.43.3) and use (III.42.2). Taking every function to be the constant one shows that the numerical constant cannot be decreased. $\square$

[write] **6 October 2026; sources 15–20.** *Credit.* For admissible slope systems this is source 06’s Theorem III.5.1, display (III.5.7); here for arbitrary surjective homomorphisms, with directional seminorms. The one-function form is Gowers’s Theorem 3.2 (`theorem_3_2`, `Sections01_03.lean`, line 258; exact companion `theorem_3_2_holds`, `Proofs03ProgressionAverage.lean`, line 381, at HEAD 0b4890a1c, for prime N and $k \leq N$).

Corollary III.43.4 (Pair-independent systems). *If $(L_i, L_j) : \Omega \rightarrow G^2$ is surjective for every $i \neq j$, then*

$$|\Lambda_L(f)| \leq \|f_i\|_{U^{k-1}} \|f_j\|_{U^{k-1}} \|f_\ell\|_2 \prod_{r \notin \{i, j, \ell\}} \|f_r\|_\infty. \quad (\text{III.43.7})$$

Proof. Pair surjectivity is equivalent to $H_{ij} = G$ for all distinct pairs, so all directional norms in Theorem III.43.3 are global norms. $\square$

Remark III.43.5 (Why two functions survive). The last step in a one-function proof bounds the penultimate function by one. In (III.43.6) we stop one step earlier and estimate the two remaining conditional means symmetrically. Their squares supply the last direction for both marked functions. This is why the output is a product, not merely the smaller of two one-function bounds.

For $k = 3$ and the usual progression forms, Fourier inversion gives a useful independent check:

$$\Lambda_3(f, g, h) = \sum_x \widehat{f}(\chi) \widehat{g}(\chi^{-2}) \widehat{h}(\chi).$$

When multiplication by two is invertible, Hölder with exponents 4, 4, 2, and Parseval, gives any appropriate two U^2 factors and one L^2 factor. Theorem III.43.3 extends this mechanism beyond the Fourier-linear case and records the exact subgroups when invertibility fails.

III.44 Sharp conversion from directional to global norms

Source 20, delivered Section 4.

Theorem III.44.1 (Finite-index comparison, with equality examples). *Let $s \geq 2$, let $H_1, \dots, H_s \leq G$, and put $m_r = [G : H_r]$. For every $f : G \rightarrow \mathbb{C}$,*

$$D_{H_1, \dots, H_s}(f) \leq \left(\prod_{r=1}^s m_r \right)^{1/2^s} \|f\|_{U^s}. \quad (\text{III.44.1})$$

For this fixed subgroup tuple, the displayed constant is optimal. In fact equality holds for every nonzero function supported on a single coset of $H_ := \bigcap_r H_r$.*

Proof. The annihilator $H_r^\perp$ consists of the m_r characters trivial on H_r . Orthogonality gives

$$m_r 1_{H_r}(h) = \sum_{\chi \in H_r^\perp} \chi(h).$$

Insert these identities into the normalized directional cube. Its 2^s -th power becomes a sum of $\prod_r m_r$ full-group cube forms, each twisted by $\prod_r \chi_r(h_r)$. The twist can be absorbed into characters multiplying the vertex functions, since $\chi_r(h_r) = \chi_r(x + h_r) \overline{\chi_r(x)}$. More explicitly, multiply the origin function by $\prod_r \overline{\chi_r}$ and the function at singleton vertex e_r by $\overline{\chi_r}$; conjugation at the singleton supplies the required positive character. All other vertex functions remain f . Lemma III.42.2 bounds the absolute value of each twisted form by $\|f\|_{U^s}^{2^s}$, using $s \geq 2$ for character invariance. Summing proves (III.44.1).

If $\text{supp } f \subset x_0 + H_*$, every nonzero cube integrand has its origin and each singleton vertex in that coset. Hence every increment lies in H_* , and in particular in its prescribed H_r . The unrestricted and restricted cube sums therefore have identical nonzero summands. Only their normalizations differ, by exactly $\prod_r [G : H_r]$. This proves equality of the powers and thus equality in (III.44.1). The all-zero list of the first $s - 1$ increments in (III.42.4) shows that $U^s(f) > 0$ for nonzero f when $s \geq 2$. Thus these examples prove optimality of the constant. $\square$

[write] **6 October 2026; sources 15–20.** *Relation to source 10.* Source 10's Theorem III.28.2 pays a kernel factor $K^{1/2^{k-1}}$ for non-invertible coefficient differences, sharp on $(\mathbb{Z}/2\mathbb{Z})^n$ (Proposition III.28.4); this theorem is the general directional-to-global conversion, with the optimal constant for every fixed subgroup tuple. With Proposition III.46.1 it answers source 06's Question 7 for the norm-conversion step (item III.9); the full multilinear optimum stays open (this source's Section V.84).

The restriction $s \geq 2$ is necessary: a nontrivial character trivial on a proper H has $D_H(f) = 1$ but $U^1(f) = 0$. Optimality in Theorem III.44.1 concerns norm conversion. It does not imply that multiplying two such constants gives the best possible constant in every multilinear counting inequality; the two equality conditions may be incompatible with the other functions.

Corollary III.44.2 (Global two-function estimate with explicit costs). *For the system (III.43.1) with $k \geq 3$, put*

$$\kappa_i = \left(\prod_{j \neq i} [G : H_{ij}] \right)^{1/2^{k-1}}. \quad (\text{III.44.2})$$

Then the right-hand side of (III.43.7), multiplied by $\kappa_i \kappa_j$, bounds $|\Lambda_L(f)|$.

Proof. Apply Theorem III.44.1 to each of the two directional norms in Theorem III.43.3. $\square$

III.44.1 Scalar progressions and composite moduli

For $L_i(x, y) = x + c_i y$, one has

$$H_{ij} = (c_i - c_j)G, \quad [G : H_{ij}] = |G[c_i - c_j]|, \quad (\text{III.44.3})$$

where $G[d] = \{x \in G : dx = 0\}$. The second identity is the finite kernel–image formula. In particular, on $\mathbb{Z}/N\mathbb{Z}$,

$$\kappa_i = \left(\prod_{j \neq i} \gcd(N, c_i - c_j) \right)^{1/2^{k-1}}. \quad (\text{III.44.4})$$

This formula remains meaningful with repeated slopes, with $\gcd(N, 0) = N$. For consecutive slopes labeled $0, \dots, k-1$,

$$\kappa_i \leq (i!(k-1-i)!)^{1/2^{k-1}}. \quad (\text{III.44.5})$$

Indeed each positive greatest common divisor is at most the corresponding nonzero integer difference. The central positions minimize this crude factorial envelope; this need not identify the minimizing positions for the actual gcd product for a given modulus.

As a small example, the slopes $0, 1, 2$ on $\mathbb{Z}/6\mathbb{Z}$ have $\kappa_0 = \kappa_2 = 2^{1/4}$ and $\kappa_1 = 1$. A bound using the middle position and either endpoint has conversion cost $2^{1/4}$; a bound using both endpoints has cost $\sqrt{2}$. Choosing the marked positions is therefore a genuinely quantitative decision when the available norm bounds are unequal.

III.44.2 Why a torsion-independent global bound can fail

Proposition III.44.3 (A collapsed-pair obstruction). *For fixed $s \geq 2$ there are one-bounded functions on groups $G_n = \mathbb{F}_2^n$ whose U^s norms tend to zero but for which a three-term progression average with two such marked functions is one. Thus no group-size-independent global two-function bound can ignore collapsed directions in arbitrary finite abelian groups.*

Proof. Assign independent uniform unit-circle phases $v(x)$ to points of a group G of size N . In the expected s -cube product, a nonzero contribution requires every point to occur equally often at even and odd vertices. In particular the even origin must agree with some odd vertex. For each of the 2^{s-1} odd vertices ω , $\omega \cdot h = 0$ has probability $1/N$, because at least one coefficient is one. A union bound gives

$$\mathbb{E}_v \|v\|_{U^s}^{2^s} \leq 2^{s-1}/N.$$

There is therefore a choice v_n with $U^s(v_n) \rightarrow 0$ on G_n . Take $f_0 = v_n$, $f_1 = 1$, $f_2 = \overline{v_n}$. In characteristic two, $x + 2y = x$, so their three-term average is $\mathbb{E}_{x,y} |v_n(x)|^2 = 1$. Both marked norms tend to zero. The subgroup $H_{02} = \{0\}$ is precisely the collapsed direction that the directional theorem retains. $\square$

III.45 Second-order telescoping and counting

Source 20, delivered Section 5.

Let $a_i : G \rightarrow [0, 1]$, let $\delta_i = \mathbb{E} a_i$, and put $b_i = a_i - \delta_i$ and $\sigma_i = \|b_i\|_2$. Then $\|a_i\|_2 \leq \sqrt{\delta_i}$ and $\sigma_i \leq \sqrt{\delta_i(1 - \delta_i)}$. For indicator functions the latter inequality is equality.

Lemma III.45.1 (First-two-residual identity). *For any scalars $a_i = \delta_i + b_i$, $1 \leq i \leq k$,*

$$\begin{aligned} \prod_{r=1}^k a_r &= \prod_{r=1}^k \delta_r + \sum_{i=1}^k b_i \prod_{r \neq i} \delta_r \\ &\quad + \sum_{1 \leq i < j \leq k} \left(\prod_{\substack{r < j \\ r \neq i}} \delta_r \right) b_i b_j \prod_{r > j} a_r. \end{aligned} \quad (\text{III.45.1})$$

Proof. Expand the product by the set of positions at which b is chosen. The empty set and singleton sets give the first two terms. Every other set has unique first two elements $i < j$. All preceding positions except i contribute δ , the positions i, j contribute b , and the remaining positions are unrestricted choices of δ or b , summing to $\prod_{r > j} a_r$. This partitions all terms exactly once. $\square$

Write

$$w_{ij} = \prod_{\substack{r < j \\ r \neq i}} \delta_r, \quad u_i = \|b_i\|_{U^{k-1}}. \quad (\text{III.45.2})$$

Surjectivity of L_i kills every singleton term after averaging. When the system is pair-independent, it also kills the term with only b_i, b_j and no tail.

Theorem III.45.2 (Quadratic counting, unequal densities). *For arbitrary surjective $L_1, \dots, L_k$ with $k \geq 3$,*

$$\left| \Lambda_L(a_1, \dots, a_k) - \prod_r \delta_r \right| \leq \sum_{i < j} w_{ij} \kappa_i \kappa_j u_i u_j. \quad (\text{III.45.3})$$

If the system is pair-independent, define, for $j < k$,

$$\mathcal{R}_j = \min \left\{ \min_{\ell > j} \sqrt{\delta_\ell}, \quad \sum_{\ell > j} \left(\prod_{j < r < \ell} \delta_r \right) \sigma_\ell \right\}. \quad (\text{III.45.4})$$

Then

$$\boxed{\left| \Lambda_L(a_1, \dots, a_k) - \prod_r \delta_r \right| \leq \sum_{i < j < k} w_{ij} u_i u_j \mathcal{R}_j.} \quad (\text{III.45.5})$$

Every permutation of the positions gives another valid bound.

Proof. Apply (III.45.1) to $a_r(L_r z)$ and average. Each summand after the singleton terms is

$$w_{ij} \mathbb{E}_z b_i(L_i z) b_j(L_j z) \prod_{r > j} a_r(L_r z). \quad (\text{III.45.6})$$

Insert constant-one functions at the unused positions. Corollary III.44.2 bounds its absolute value by $w_{ij} \kappa_i \kappa_j u_i u_j$, since all unmarked factors are one-bounded. Summing proves (III.45.3).

Now assume pair independence. When $j = k$, the average in (III.45.6) factors as $(\mathbb{E} b_i)(\mathbb{E} b_j) = 0$. For $j < k$, choose any $\ell > j$ as the L^2 factor in Corollary III.43.4. This gives the first bound $u_i u_j \min_{\ell > j} \sqrt{\delta_\ell}$ for the average in (III.45.6).

For the second bound use the exact first-residual identity in the tail:

$$\prod_{r>j} a_r = \prod_{r>j} \delta_r + \sum_{\ell>j} \left(\prod_{j<r<\ell} \delta_r \right) b_\ell \prod_{r>\ell} a_r.$$

The constant-tail contribution vanishes by pair independence. In each remaining contribution use b_ℓ as the L^2 factor, retaining b_i, b_j as marked functions. This gives the second expression in (III.45.4). Taking the smaller bound for each (i, j) and summing proves (III.45.5). The proof is unchanged by a permutation of the positions. $\square$

Remark III.45.3 (Order-sensitive norms). In (III.45.5), $u_i u_j$ can be replaced term by term by $\|b_i\|_{U^{k-j+1}} \|b_j\|_{U^{k-j+1}}$. Indeed the active subsystem in (III.45.6) has $k - j + 2$ positions, and the same proof applies to it. The tail expansion does not increase that number of active positions. Equivalently one may retain constant-one positions within that subsystem. The uniform highest-order version is simpler to state; the order-sensitive version can be useful in a formal or numerical optimization of the bound.

III.45.1 A single set and an explicit coefficient

Corollary III.45.4 (Symmetric quadratic counting). *Let $A \subset G$ have density $0 < \delta < 1$, and assume pair independence of the k forms. Put $u = \|1_A - \delta\|_{U^{k-1}}$. Then*

$$\left| \Lambda_L(1_A, \dots, 1_A) - \delta^k \right| \leq B_k(\delta) u^2,$$

with B_k as in (F.25.3). Equations (F.25.5) and (F.25.6) hold.

Proof. There are $j - 1$ choices of $i < j$, each with $w_{ij} = \delta^{j-2}$. The two tail bounds in (III.45.4) become $\sqrt{\delta}$ and $\sigma \sum_{r=0}^{k-j-1} \delta^r$. This proves the formula for B_k . For $k = 3$ only $j = 2$ occurs, and the smaller tail bound is σ . Using $R_m \leq \sqrt{\delta}$ and the positive series $\sum_{r \geq 0} (r+1) \delta^r = (1 - \delta)^{-2}$ proves (F.25.5).

For $0 < \delta \leq 1/2$ and $m \geq 2$,

$$\frac{\sigma(1 + \delta + \dots + \delta^{m-1})}{\sqrt{\delta}} \geq \sqrt{1 - \delta}(1 + \delta) \geq 1.$$

The last inequality follows on squaring from $\delta(1 - \delta - \delta^2) \geq 0$ in this range. Thus $R_m = \sqrt{\delta}$ for $m \geq 2$, whereas $R_1 = \sigma$. Separating the last term of the sum gives (F.25.6). $\square$

[write] **6 October 2026; sources 15–20.** *Against source 06* (proved at the write). Source 06's Corollary III.5.5 has the coefficient $\sqrt{\delta} B_k^{(06)}(\delta)$ with $B_k^{(06)}(\delta) = \sum_{t=0}^{k-3} (t+1) \delta^t$ (III.5.10). Since $R_m(\delta) \leq \sqrt{\delta}$, putting $t = j - 2$ in (F.25.3) gives $B_k(\delta) \leq \sqrt{\delta} \sum_{t=0}^{k-3} (t+1) \delta^t = \sqrt{\delta} B_k^{(06)}(\delta)$. The inequality is strict for every $k \geq 3$ and $0 < \delta < 1$: the term $j = k - 1$ carries $R_1(\delta) = \sigma < \sqrt{\delta}$ with the positive weight $(k - 2) \delta^{k-3}$. For $k = 3$ the coefficient is σ , source 06's sharp Proposition III.5.6. Ratios: 0.8536 at $k = 4$, $\delta = 1/2$; $\sqrt{0.7} = 0.8367$ at $k = 3$, $\delta = 0.3$. This advances source 06's Question 6 without settling it (item III.8).

This coefficient is the best obtained by the two explicitly specified tail estimates at each term of this fixed telescoping identity. No claim is made that it is the optimal coefficient among all counting arguments.

δ	k	$S_{k-2}(\delta)$	$B_k(\delta)$	$\sqrt{\delta^k/(2B_k(\delta))}$
0.10	3	1.000000	0.300000	0.04082483
0.10	5	1.110000	0.388473	0.00358760
0.10	8	1.111110	0.390401	0.00011317
0.25	3	1.000000	0.433013	0.13432124
0.25	5	1.312500	0.831190	0.02423734
0.25	8	1.333008	0.887303	0.00293230
0.50	3	1.000000	0.500000	0.35355339
0.50	5	1.750000	1.789214	0.09344992
0.50	8	1.968750	2.612818	0.02734075

Table III.45.1: Computed coefficients. The old and new error bounds are respectively $S_{k-2}u$ and $B_k u^2$, not $S_{k-2}u^2$ and $B_k u^2$. The last column is the sufficient uniformity threshold used below.

III.45.2 A torsion-inclusive symmetric statement

Without pair independence, Corollary III.44.2 and (III.45.3) give

$$\left| \Lambda_L(1_A, \dots, 1_A) - \delta^k \right| \leq C_{L,\delta} u^2, \quad C_{L,\delta} = \sum_{j=2}^k \delta^{j-2} \kappa_j \sum_{i < j} \kappa_i. \quad (\text{III.45.7})$$

In particular $C_{L,\delta} \leq \kappa_{\max}^2/(1-\delta)^2$. This is an explicit bound on every finite abelian group. It may be weak when a direction collapses, and Proposition III.44.3 explains why some loss is then unavoidable.

III.46 Pair correlations as an exact torsion correction

Source 20, delivered Section 6.

The directional estimate can be complemented by an exact Fourier description of the first obstruction to a product-density baseline.

Proposition III.46.1 (Pair correction and its sharp constant). *Let $f, g : G \rightarrow \mathbb{C}$ have mean zero, let d be an integer, and put $H = dG$, $m = [G : H]$. Then*

$$Q_d(f, g) := \mathbb{E}_{x,y} f(x)g(x+dy) = \sum_{\substack{\chi \in H^\perp \\ \chi \neq 1}} \hat{f}(\chi) \hat{g}(\chi^{-1}), \quad (\text{III.46.1})$$

$$|Q_d(f, g)| \leq \sqrt{m-1} \|f\|_{U^2} \|g\|_{U^2}. \quad (\text{III.46.2})$$

For $m > 1$, the constant $\sqrt{m-1}$ is optimal, also when the functions are required to be real and one-bounded.

Proof. Fourier inversion and averaging x force the two frequencies to be inverses. Averaging dy , uniform in H , forces the surviving character to lie in $H^\perp$. The zero-frequency term vanishes by centering. Hölder on the remaining $m-1$ frequencies, with exponents 4, 4, 2, and (III.42.5) prove (III.46.2). For sharpness take

$$f = g = t \sum_{\substack{\chi \in H^\perp \\ \chi \neq 1}} \chi = t(m1_H - 1), \quad 0 < t \leq 1/(m-1).$$

These are real, centered and one-bounded. Their pair correlation is $(m-1)t^2$, while each fourth power of the U^2 norm is $(m-1)t^4$. The ratio is exactly $\sqrt{m-1}$. $\square$

For $a_i = \delta_i + b_i$, the full balanced expansion separates an exact pair contribution

$$\sum_{i < j} \left(\prod_{r \notin \{i, j\}} \delta_r \right) Q_{c_j - c_i}(b_i, b_j) \quad (\text{III.46.3})$$

from terms with at least three balanced factors. Thus on a non-admissible cyclic group, keeping (III.46.3) can be preferable to regarding it as error. In the admissible case every summand is zero. The finite-index factors and this quotient-frequency correction describe the same obstruction in different languages: one is suited to Cauchy–Schwarz iteration, the other isolates the precise pair bias.

III.47 Supersaturation and the uniformity parameter

Source 20, delivered Section 7.

A *proper ordered progression* is a pair (x, y) whose k entries $x, x + y, \dots, x + (k-1)y$ are distinct and belong to A . Let $T_k^*(A)$ denote their number.

Corollary III.47.1 (Admissible supersaturation). *Suppose the consecutive progression system is admissible, $k \geq 3$, $|G| = N$, and $A \subset G$ has density $0 < \delta < 1$. With $u = \|1_A - \delta\|_{U^{k-1}}$,*

$$T_k^*(A) \geq N^2(\delta^k - B_k(\delta)u^2) - \delta N. \quad (\text{III.47.1})$$

In particular $T_k^(A) > 0$ whenever*

$$u^2 \leq \frac{\delta^k}{2B_k(\delta)}, \quad N > 2\delta^{1-k}. \quad (\text{III.47.2})$$

Proof. The normalized count includes all N^2 parameter pairs. A repeated entry implies $(r-s)y = 0$ for $0 < |r-s| < k$, hence $y = 0$ by admissibility. Exactly $|A| = \delta N$ counted pairs have zero difference. Subtract these from the lower bound supplied by Corollary III.45.4. Conditions (III.47.2) make the resulting lower bound strictly positive. $\square$

[write] **6 October 2026; sources 15–20.** *Comparison* (at the write). Source 06’s Corollary III.5.7 needs $u^2 \leq \delta^{k-1/2}/(2B_k^{(06)}(\delta)) = \delta^k/(2\sqrt{\delta}B_k^{(06)}(\delta))$; since $B_k(\delta) \leq \sqrt{\delta}B_k^{(06)}(\delta)$, the condition (III.47.2) is weaker, so this criterion is the stronger one; for the other criteria of this report see the note after Corollary III.4.4.

In Gowers’s degree- $k-2$ convention, the power parameter corresponding to the normalized cube norm is

$$\alpha = u^{2^{k-1}}. \quad (\text{III.47.3})$$

Thus the new counting error is $B_k(\delta)\alpha^{1/2^{k-2}}$, and the sufficient condition becomes

$$\alpha \leq \left(\frac{\delta^k}{2B_k(\delta)} \right)^{2^{k-2}}. \quad (\text{III.47.4})$$

For $0 < \delta \leq 1/2$, the convenient bound $B_k(\delta) \leq 4\sqrt{\delta}$ gives the simpler sufficient condition

$$\alpha \leq \left(\frac{\delta^{k-1/2}}{8} \right)^{2^{k-2}}. \quad (\text{III.47.5})$$

This improves the local symmetric threshold derived from (F.25.1), which is $\alpha \leq (\delta^k / (2S_{k-2}(\delta)))^{2^{k-1}}$. The comparison keeps the same normalized u and distinguishes a cube norm from its 2^{k-1} -th power.

III.47.1 Removing degenerate progressions with torsion

For an arbitrary finite abelian group put

$$D_k(G) = \left| \bigcup_{r=1}^{k-1} G[r] \right|.$$

The difference y gives a repeated entry precisely when its order is less than k . For each such difference at most $|A|$ starting points can contribute. Therefore (III.45.7) gives

$$T_k^*(A) \geq N^2(\delta^k - C_{L,\delta}u^2) - \delta N D_k(G). \quad (\text{III.47.6})$$

For cyclic groups this degeneracy count has the exact formula

$$D_k(\mathbb{Z}/N\mathbb{Z}) = \sum_{\substack{d|N \\ d < k}} \varphi(d), \quad (\text{III.47.7})$$

since there are $\varphi(d)$ elements of order d . This includes the zero difference through $d = 1$.

III.47.2 The no-wrap distinction

Lemma III.47.2 (A sufficient lifting condition). *Suppose $A \subset \{0, \dots, n-1\}$ is embedded in $\mathbb{Z}/N\mathbb{Z}$ and $2(n-1) < N$. Any nonconstant modular progression of length at least three contained in A is an ordinary integer progression; reversing its order if necessary gives positive difference.*

Proof. Let $a_0, \dots, a_{k-1}$ be its integer representatives. Their consecutive differences are congruent modulo N . Each lies between $-(n-1)$ and $n-1$, so the difference between any two has absolute value less than N and must be zero. They are all the same integer, which is nonzero for a nonconstant progression. $\square$

The lemma does not show that cutting an arbitrary uniform set down to a short interval preserves its balanced norm. Such a localization step can introduce unbalanced cutoff functions at precisely the positions where our quadratic argument needs small norms. This is one reason (III.47.4) is not silently promoted to a new global Szemerédi estimate. The inverse theorem, density increment, iteration losses, and localization must all be analyzed together.

III.48 A sharp quadratic obstruction for three-term counts

Source 20, delivered Section 8.

We now prove that quadratic dependence in (F.25.4) is the correct uniform exponent when $k = 3$, not just an artifact of stopping Cauchy–Schwarz early. The construction is explicit for bounded weights; a second argument transfers its scale to actual subsets.

III.48.1 An elementary even-dimensional Gauss identity

Let $p = 29$, let $G = \mathbb{F}_p^{2m}$, $N = p^{2m}$, and set $Q(x) = x_1^2 + \cdots + x_{2m}^2$. Write $e_p(t) = \exp(2\pi it/p)$.

Lemma III.48.1 (Gauss averages). *For every nonzero $t \in \mathbb{F}_p$,*

$$\mathbb{E}_{x \in G} e_p(tQ(x)) = N^{-1/2}, \quad \left| \widehat{e_p(tQ)}(\xi) \right| = N^{-1/2} \quad (\xi \in \widehat{G}). \quad (\text{III.48.1})$$

More generally, if a quadratic form in two scalar variables over $\mathbb{F}_p$ has rank r , its sum over $2m$ independent coordinate copies has normalized exponential average $N^{-r/2}$.

Proof. Let $g(t) = \sum_{z \in \mathbb{F}_p} e_p(tz^2)$. If χ is the quadratic character, counting square roots gives $g(t) = \chi(t)g(1)$ for $t \neq 0$. The change of variables $(z, w) \mapsto (z - w, z + w)$ is invertible and gives $|g(t)|^2 = \sum_{z, w} e_p(t(z^2 - w^2)) = p$. Also $\overline{g(1)} = g(-1) = \chi(-1)g(1) = g(1)$ because $p \equiv 1$ modulo four. Hence $g(t)^2 = p$, irrespective of its possible sign. The product over $2m$ coordinates is p^m , and normalization gives the first assertion of (III.48.1). Completing the square for each linear Fourier twist proves the magnitude assertion.

Every symmetric quadratic matrix over this odd-characteristic field is diagonalizable by an invertible linear change of variables. A rank- r form has r nonzero diagonal coefficients; zero directions contribute one. For each nonzero coefficient the $2m$ coordinate copies contribute $p^{-m} = N^{-1/2}$. Their product is $N^{-r/2}$. $\square$

III.48.2 Four phases and an exact rank census

Set $a = 1/16$ and

$$h(x) = a \sum_{t \in \{-2, -1, 1, 2\}} e_p(tQ(x)), \quad f(x) = h(x) - 4aN^{-1/2}, \quad w(x) = \frac{1}{2} + f(x). \quad (\text{III.48.2})$$

The conjugate pairs make h, f, w real. Their means are $\mathbb{E}h = 4aN^{-1/2}$, $\mathbb{E}f = 0$, and $\mathbb{E}w = 1/2$. Moreover $|f| \leq 4a(1 + N^{-1/2}) \leq 1/2$, so $0 \leq w \leq 1$.

Proposition III.48.2 (Exact weighted excess). *For the functions in (III.48.2),*

$$\|f\|_{U^2}^4 \leq (4a)^4/N, \quad (\text{III.48.3})$$

$$\Lambda_3(w) - \frac{1}{8} = a^3 \left(2N^{-1/2} + 62N^{-1} - 64N^{-3/2} \right). \quad (\text{III.48.4})$$

Consequently, as $m \rightarrow \infty$, $\|f\|_{U^2} \rightarrow 0$ and

$$\Lambda_3(w) - \frac{1}{8} \geq \frac{1}{128} \|f\|_{U^2}^2. \quad (\text{III.48.5})$$

Proof. Centering removes the zero Fourier coefficient. Every other coefficient of f has absolute value at most $4aN^{-1/2}$, by Lemma III.48.1. There are at most N coefficients, so (III.42.5) proves (III.48.3).

Expand $\Lambda_3(h)$ by a triple $(t_0, t_1, t_2) \in \{-2, -1, 1, 2\}^3$. At each coordinate, the quadratic form in (x, y) has symmetric matrix

$$M_t = \begin{pmatrix} t_0 + t_1 + t_2 & t_1 + 2t_2 \\ t_1 + 2t_2 & t_1 + 4t_2 \end{pmatrix}, \quad \det M_t = t_0t_1 + 4t_0t_2 + t_1t_2. \quad (\text{III.48.6})$$

Its determinant has absolute value at most 24, so vanishing modulo 29 is equivalent to integer vanishing. The only vanishing triples are $(-1, 2, -1)$ and $(1, -2, 1)$. To check this without a search oracle, solve

$$t_1(t_0 + t_2) = -4t_0t_2.$$

If $t_0 + t_2 = 0$, the right side is nonzero. For equal endpoints the solution is $t_1 = -2t_0$, which is allowed exactly for $t_0 = \pm 1$. For unequal endpoints of the same sign the required middle value has magnitude $8/3$; for unequal endpoints of opposite signs it has magnitude 8. Neither is allowed. These exhaust the possible endpoint pairs. No matrix has rank zero: subtracting its off-diagonal entry from the bottom-right entry would give $2t_2 = 0$. Thus two of the 64 matrices have rank one and 62 have rank two. Lemma III.48.1 gives

$$\Lambda_3(h) = a^3(2N^{-1/2} + 62N^{-1}).$$

Any two progression positions are independent because p is odd. Expanding the centered triple therefore gives $\Lambda_3(f) = \Lambda_3(h) - (\mathbb{E}h)^3$. The singleton and pair terms in $\Lambda_3(1/2 + f)$ vanish. This proves (III.48.4). For $N \geq 2$, its right-hand side is at least $2a^3N^{-1/2}$, whereas (III.48.3) implies $\|f\|_{U^2}^2 \leq 16a^2N^{-1/2}$. Their ratio is at least $a/8 = 1/128$. $\square$

This construction supplies small global U^2 norm through the large rank of the quadratic phase, while the two rank-one combinations in (III.48.6) preserve a larger three-term correlation. Scaling a fixed cosine alone would not establish the same exponent: it would give a cubic correlation in its small amplitude.

III.48.3 Rounding without losing the critical scale

A naive assertion that a random rounding is “close” is insufficient: the excess itself is of order $N^{-1/2}$. We need simultaneous control of a positive count excess and a fourth power of the norm of order N^{-1} . The following argument supplies that control explicitly.

Theorem III.48.3 (Quadratic sharpness for indicator sets). *There are subsets $A_m \subset \mathbb{F}_{29}^{2m}$, for all sufficiently large m , with densities $\delta_m \rightarrow 1/2$ and $u_m = \|1_{A_m} - \delta_m\|_{U^2} \rightarrow 0$, such that*

$$\Lambda_3(1_{A_m}) - \delta_m^3 \geq c_* u_m^2 \tag{III.48.7}$$

for an absolute $c_* > 0$. Consequently, for every $\eta > 0$, no constant C_η can bound all these errors by $C_\eta u_m^{2+\eta}$. The failure also holds uniformly for densities in any fixed neighborhood of $1/2$.

Proof. Include each x independently with probability $w(x)$ from (III.48.2). Let $B_x = 1_A(x)$, $\delta_A = \mathbb{E}_x B_x$, and put

$$X_A = \Lambda_3(1_A) - \delta_A^3, \quad S_A = \|1_A - \delta_A\|_{U^2}^4.$$

We first show, for all sufficiently large N , that

$$\mathbb{E}_A X_A \geq cN^{-1/2}, \quad \mathbb{E}_A S_A \leq C/N, \quad c = a^3, \quad C = 8((4a)^4 + 2) = 513/32. \tag{III.48.8}$$

All three points of a progression are distinct unless $y = 0$. Off that diagonal, independence gives the product of the three probabilities. On the diagonal, $B_x^3 = B_x$ and $w(x) \geq w(x)^3$. Therefore $\mathbb{E}_A \Lambda_3(1_A) \geq \Lambda_3(w)$. Writing $\xi_x = B_x - w(x)$ and expanding the cube of their average gives

$$\mathbb{E}_A \delta_A^3 = \delta^3 + \frac{3\delta}{N} \mathbb{E}_x w(x)(1 - w(x)) + \frac{1}{N^2} \mathbb{E}_x w(x)(1 - w(x))(1 - 2w(x)), \tag{III.48.9}$$

where $\delta = 1/2$. Since $w(1-w) \leq 1/4$, Proposition III.48.2 yields

$$\mathbb{E}_A X_A \geq 2a^3 N^{-1/2} - \frac{3}{8N} - \frac{1}{4N^2} \geq a^3 N^{-1/2}$$

for sufficiently large N .

For the second assertion, note that $1_A - \delta_A = f + (\xi - \mathbb{E}_x \xi_x)$. By (III.42.5), subtracting the mean removes a nonnegative term from the fourth power of a U^2 norm. Also

$$\mathbb{E}_A \|\xi\|_{U^2}^4 \leq 2/N. \quad (\text{III.48.10})$$

To prove (III.48.10), expand the parallelogram. Independence and $\mathbb{E}_A \xi_x = 0$ kill any term with a point appearing exactly once. In an odd-order group, a parallelogram with no singleton among its four vertices must have $h = 0$ or $k = 0$. Indeed its three possible pairings give respectively $h = 0$, $k = 0$, or $h = k$ and $2h = 0$; the last case is also zero. The union has normalized measure at most $2/N$, and the absolute value of each expected product is at most one. The triangle inequality and $(v+w)^4 \leq 8(v^4 + w^4)$ now give

$$\mathbb{E}_A S_A \leq 8 \left(\|f\|_{U^2}^4 + \mathbb{E}_A \|\xi\|_{U^2}^4 \right) \leq C/N,$$

which proves (III.48.8).

The already proved three-term counting inequality implies $|X_A| \leq \sqrt{S_A}$ for every realization (in fact the prefactor is at most $1/2$). Let $E = \{X_A \geq c/(2\sqrt{N})\}$. By Cauchy–Schwarz,

$$\frac{c}{\sqrt{N}} \leq \mathbb{E}_A X_A \leq \frac{c}{2\sqrt{N}} + \mathbb{E}_A(1_E \sqrt{S_A}) \leq \frac{c}{2\sqrt{N}} + \sqrt{C/N} \sqrt{\Pr(E)}.$$

Consequently $\Pr(E) \geq c^2/(4C)$. Set $K = 8C^2/c^2$. Markov's inequality gives $\Pr(S_A > K/N) \leq c^2/(8C)$. Finally $\text{Var}(\delta_A) \leq 1/(4N)$, so

$$\Pr(|\delta_A - 1/2| > N^{-1/4}) \leq 1/(4\sqrt{N}).$$

For all sufficiently large N , these estimates leave a positive probability of all three events

$$X_A \geq c/(2\sqrt{N}), \quad S_A \leq K/N, \quad |\delta_A - 1/2| \leq N^{-1/4}.$$

Choose one such realization for each large m . Then $u_m^2 = \sqrt{S_{A_m}} \leq \sqrt{K/N}$ and (III.48.7) holds with $c_* = c/(2\sqrt{K}) > 0$. The norm cannot be zero when $X_A > 0$. Dividing by $u_m^{2+\eta}$ now yields a lower bound $c_* u_m^{-\eta} \rightarrow \infty$, proving the claimed impossibility. $\square$

[write] 6 October 2026; sources 15–20. New in this report: source 06's three-term coefficient is sharp only at singletons ($\delta = 1/N$); this theorem shows that the exponent 2 cannot be raised at fixed density near $1/2$, for indicator sets on $\mathbb{F}_{29}^{2m}$. Not for prime cyclic groups or $k \geq 4$ (the source's non-claims; item III.8).

The constant c_* is intentionally crude. The theorem determines an exponent, not the optimal coefficient in the indicator-set extremal problem. The finite check of the weight construction uses $m = 1$; the existence proof for the rounded sets is asymptotic and does not claim that its crude probability thresholds are useful computationally.

III.49 Source 21's chapter: the exponent-five coefficient, with source 23's second route [write]

Added in the fifth write, 6 October 2026.

The next ten sections settle the question that source 14's Theorem III.35.2 left open: on groups of exponent five, the first correction after $6\sqrt{2}\delta^4u^4$ in the centered three-cube problem has exponent six, and its coefficient is now known. Sections III.50–III.57 are source 21, the main source; Sections III.58–III.59 are source 23's Section 7 and Appendix A, an independent proof of the same theorem, written the same day from a pin three minutes later, printed here as a credited second route. Neither source cites the other; both cite source 14's manuscript. *Logically the chapter follows source 14's Section III.35*; it is printed here, at the end of Part III, because no number printed earlier may move. Source 21's introduction, with the main theorem (Theorem F.26.1), is Section F.26; its repository interface, questions and appendices are in Part V (Sections V.87–V.90). Source 23's introduction is Section F.28. Their notation is in the table at the head of this Part.

What is proved, in one paragraph. For real centered f on any nontrivial finite abelian group of exponent five and $u = \|f\|_{U^3} \leq c\delta$, $Q_3(\delta + f) \leq \delta^8 + 6\sqrt{2}\delta^4u^4 + C_5\delta^2u^6 + Cu^8$ with $C_5 = 2^{3/4}/3$ and absolute c, C , sharp on every fixed such group (source 21's Theorem F.26.1, Corollary III.55.2; source 23's Theorem III.58.1). The mechanism is the same in both: a positive cubic statistic controls the sixth- and seventh-degree terms through the Fourier tail, and the order-five resonance makes the cost of the second harmonic quadratic. Source 21 also proves the phase envelope $C_5(\cos \theta)_+^2/(1 + \cos^2 \phi)$ (Theorem III.53.2), rigidity of near-extremizers (Theorem III.54.1; source 23's Theorem III.58.3 is equivalent), the next coefficient 20/9 on $\mathbb{Z}/5\mathbb{Z}$ for the maximum (Theorem III.55.1; source 23's Appendix A computes 20/9 for one model, a lower bound), indicator sharpness through u^8 as lower examples (Theorem III.56.1; source 23 through u^6) and a projection consequence of cube Cauchy–Schwarz (Proposition III.57.1). The universal coefficient of u^8 is open (item III.27).

Relation to the formal statements (checked at HEAD 2dfe46f44). `lemma_3_10` (Sections01_0 3.lean, line 340) is an exact companion (`lemma_3_10_holds`, `Proofs03CubeUpper.lean`, line 82), as source 21 records; no characteristic-five or cube-expansion declaration exists. Nothing in this chapter is formalized.

III.50 Fourier identities and the sharp quartic deficit

Source 21, delivered Section 2.

For $f : G \rightarrow \mathbb{C}$, our Fourier convention is

$$a_r = \widehat{f}(r) = \mathbb{E}_{x \in G} f(x) \overline{r(x)}, \quad f(x) = \sum_{r \in \widehat{G}} a_r r(x).$$

Throughout the main argument f is real and centered, so $a_{-r} = \overline{a_r}$ and $a_0 = 0$. Set

$$\lambda_r = |a_r|^2, \quad S = \sum_r |a_r|^4 = \|f\|_{U^2}^4, \quad M = \sum_r |a_r|^8. \quad (\text{III.50.1})$$

All sums on the dual group are unnormalized.

Lemma III.50.1 (A sum of Fourier squares). *Define*

$$T(s, t) = \sum_r a_r \overline{a_{r+s}} \overline{a_{r+t}} a_{r+s+t}. \quad (\text{III.50.2})$$

Then

$$\|f\|_{U^3}^8 = \sum_{s,t} |T(s,t)|^2. \quad (\text{III.50.3})$$

For real f , $T(s,t)$ is real, is symmetric in s,t , and is unchanged by replacing either s or t by its negative.

Proof. Insert Fourier inversion in the eight-vertex cube. Averaging over x, h_1, h_2, h_3 imposes the four equations

$$\sum_{\omega} r_{\omega} = 0, \quad \sum_{\omega} \omega_j r_{\omega} = 0 \quad (1 \leq j \leq 3).$$

Their solutions are precisely the parity-adjusted affine frequency cubes. Grouping two opposite faces gives the expansion of the right side of (III.50.3). Equivalently, one can expand the right side and use its four free frequencies as parameters for the same four equations. This parameterization uses only invertible integer changes of variables, so the identity holds on every finite abelian group.

In (III.50.2), the substitution $r \mapsto -r - s - t$ pairs each summand with its complex conjugate because $a_{-r} = \overline{a_r}$. Hence T is real. Interchanging s, t is immediate; shifting the base r after changing the sign of one increment conjugates T , and so leaves it unchanged in the real case. $\square$

Define the positive energy

$$\mathcal{E}(\lambda) = \sum_{r+s=t+w} \lambda_r \lambda_s \lambda_t \lambda_w = \sum_v (\lambda * \lambda)(v)^2,$$

where the convolution is an unnormalized sum. Since λ is even, $T(0,t) = \sum_r \lambda_r \lambda_{r+t}$, and $T(0,0) = S$. The two axes of (III.50.3) therefore contribute exactly

$$2\mathcal{E}(\lambda) - S^2. \quad (\text{III.50.4})$$

Lemma III.50.2 (Quartic comparison and concentration). *For every real centered f on an odd-order finite abelian group,*

$$u^8 \geq 5S^2 - 6M \geq 2S^2, \quad u = \|f\|_{U^3}. \quad (\text{III.50.5})$$

If $f \neq 0$, choose a largest conjugate pair $\{\xi, -\xi\}$ and put

$$\rho = |a_{\xi}|, \quad D = \sum_{r \notin \{\xi, -\xi\}} |a_r|^4, \quad S = 2\rho^4 + D. \quad (\text{III.50.6})$$

Then

$$u^8 \geq 2S^2 + 3DS. \quad (\text{III.50.7})$$

In particular, equality in $u^8 \geq 2S^2$ for nonzero f forces exactly one conjugate Fourier pair. In exponent five, every such pair gives equality.

Proof. Write the energy as a sum over $r + s + t + w = 0$. Retain the union of the three families obtained by pairing the four entries into opposite pairs. Each family has weight S^2 , every pairwise intersection has weight M , and the triple intersection has weight zero: it requires $2r = 0$, and the zero Fourier coefficient vanishes. Thus $\mathcal{E}(\lambda) \geq 3S^2 - 3M$. Combine this with (III.50.4).

The nonzero Fourier coefficients come in conjugate pairs. If their fourth-power masses are $x_j = 2|a_{r_j}|^4$, then $S = \sum_j x_j$ and $M = \frac{1}{2} \sum_j x_j^2 \leq S^2/2$. This proves the second inequality and its necessary equality case. Maximality also gives $M \leq \rho^4 S = (S - D)S/2$, yielding (III.50.7). For a pure pair of order five, directly evaluating (III.50.3) gives $u^8 = 8\rho^8 = 2S^2$. $\square$

[write] **6 October 2026; sources 21–27.** *Credit.* The inequality $u^8 \geq 2S^2$ for real centered f on odd-order groups is source 04's Theorem III.3.4; source 14 reproves it (its display (III.34.8)). The refinement (III.50.7) with the tail D is the form used below.

The last sufficiency statement was restricted deliberately. For example, short resonances at order three change the norm of a pure pair. The present report does not silently extend an exponent-five equality formula to all odd orders.

III.51 Higher centered coefficients with a retained Fourier tail

Source 21, delivered Section 3.

The following exact expansion introduces the terms whose crude treatment prevented the earlier coefficient from being identified.

Lemma III.51.1 (Centered cube expansion). *On an exponent-five group,*

$$Q_3(\delta + f) = \delta^8 + 12\delta^4 S + 8\delta^3 B_5(f) + \delta^2 C_6(f) + \delta C_7(f) + u^8, \quad (\text{III.51.1})$$

where C_j is the sum of all j -vertex centered terms and

$$B_5(f) = \sum_r a_{2r} a_{-r}^3 a_r. \quad (\text{III.51.2})$$

The number $B_5(f)$ is real.

Proof. Expand the product over subsets of the eight vertices. One, two, or three distinct cube forms are jointly uniform, so centering kills their expectations. Among the 70 four-element supports, the 12 parallelograms have a one-dimensional relation and contribute S ; the other 58 supports are independent over $\mathbb{F}_5$.

Of the 56 five-element supports, 48 have a relation with a zero coefficient, forcing a factor a_0 in their Fourier expectation. The remaining eight are cube-symmetric copies of $\{000, 100, 010, 001, 111\}$. If the frequency at 111 is r , the four orthogonality equations force frequencies $2r, -r, -r, -r, r$ on this support. This proves (III.51.2). Negating r conjugates a summand, so the sum is real.

The support counts may also be checked by row reduction of the columns $(1, \omega_1, \omega_2, \omega_3)$, as described in Appendix V.89. The nonzero determinants needed for these classifications have absolute value at most two; in characteristic five their ranks agree with the stated ranks. $\square$

Cube Cauchy–Schwarz alone gives $|C_6| \leq 28u^6$ and $|C_7| \leq 8u^7$. The first of these bounds is too large to determine the coefficient of u^6 . We instead use the positive statistic

$$H = \sum_{r+s+t=0} \lambda_r \lambda_s \lambda_t. \quad (\text{III.51.3})$$

Theorem III.51.2 (Sextic and septic profile inequalities). *For every real centered f on an exponent-five group,*

$$|C_6(f)| \leq 28H, \quad |C_7(f)| \leq 8u^4 \sqrt{H}. \quad (\text{III.51.4})$$

For a largest pair and tail D as in (III.50.6),

$$H \leq 3\sqrt{D\mathcal{E}(\lambda)} \leq \frac{3\sqrt{3}}{2} u^4 \sqrt{D}. \quad (\text{III.51.5})$$

Consequently,

$$\begin{aligned} |C_6(f)| &\leq 42\sqrt{3}u^4D^{1/2}, \\ |C_7(f)| &\leq 8\sqrt{3\sqrt{3}/2}u^6D^{1/4}. \end{aligned} \tag{III.51.6}$$

No bound on $\|f\|_\infty$ is assumed.

Proof. Classify a six-vertex support by the two omitted vertices. Their Hamming distance is one, two, or three. The respective numbers of unordered pairs are 12, 12, 4. Fourier orthogonality gives the following representatives, with all frequency sums over $\hat{G}$:

Omitted vertices	Multiplicity	Six-vertex expectation
000, 001	12	H
000, 011	12	$J = \sum_{r,s} \lambda_r a_s^2 a_{r-s} \overline{a_{r+s}}$
000, 111	4	$K = \sum_{r+s+t=0} a_r^2 a_s^2 a_t^2$

Each formula follows by solving the same four equations used in Lemma III.50.1, now with the omitted frequencies set to zero. Clearly $|K| \leq H$. For J , Cauchy–Schwarz gives

$$\begin{aligned} |J| &\leq \sum_{r,s} \lambda_r \lambda_s \sqrt{\lambda_{r-s} \lambda_{r+s}} \\ &\leq \left(\sum_{r,s} \lambda_r \lambda_s \lambda_{r-s} \right)^{1/2} \left(\sum_{r,s} \lambda_r \lambda_s \lambda_{r+s} \right)^{1/2} = H, \end{aligned}$$

using evenness of λ . This proves the sextic bound.

For the septic bound, replace the zero Fourier coefficient in T by a real variable δ and differentiate at $\delta = 0$. The derivative is

$$L(s, t) = 2 \operatorname{Re}(a_s a_t \overline{a_{s+t}}) + 2 \operatorname{Re}(a_s \overline{a_t} \overline{a_{s-t}}). \tag{III.51.7}$$

This is the sum of the four terms obtained by choosing which of the four factors contains the differentiated zero coefficient. If indices coincide, the extra zero coefficients make the corresponding terms vanish, so the formula still applies. Each of these four cubic terms has squared ℓ^2 norm H . Thus $\|L\|_2 \leq 4\sqrt{H}$. Differentiating the sum of squares gives

$$C_7(f) = 2 \operatorname{Re}\langle T, L \rangle, \quad |C_7(f)| \leq 2\|T\|_2 \|L\|_2 \leq 8u^4\sqrt{H}.$$

Finally, no triple drawn from $\{\xi, -\xi\}$ sums to zero because ξ has order five. Every nonzero summand in H contains at least one tail index. Therefore

$$H \leq 3 \sum_{r \notin \{\xi, -\xi\}} \lambda_r (\lambda * \lambda)(-r) \leq 3\sqrt{D\mathcal{E}(\lambda)}.$$

The axes in (III.50.3) give $\mathcal{E}(\lambda) \leq (u^8 + S^2)/2$. Lemma III.50.2 gives $S^2 \leq u^8/2$, proving (III.51.5) and (III.51.6). $\square$

Remark III.51.3 (Why the two improvements must be used together). The new bounds express the higher coefficients in terms of the same tail D that lowers the quartic contribution. A merely qualitative statement that C_6 tends to zero near a pure pair would give a sharp coefficient with an unspecified little-oh error. The powers $D^{1/2}$ and $D^{1/4}$ allow the explicit $O(u^8)$ remainder.

III.52 The characteristic-five resonance and a finite bound

Source 21, delivered Section 4.

For a largest pair, set $b = |a_{2\xi}|$. The pair $\{2\xi, -2\xi\}$ is distinct from $\{\xi, -\xi\}$, so $2b^4 \leq D$. When $b > 0$, write

$$a_\xi = \rho e^{i\alpha}, \quad a_{2\xi} = be^{i\beta}, \quad \theta = 2\alpha - \beta, \quad \phi = 3\alpha + \beta. \quad (\text{III.52.1})$$

These phases are understood modulo 2π . Notice the useful identity $\theta + \phi = 5\alpha$.

Theorem III.52.1 (A joint norm penalty). *With the preceding notation,*

$$u^8 \geq 2S^2 + 12\rho^4 D + 16\rho^6 b^2. \quad (\text{III.52.2})$$

Let $D_{\text{out}} = \sum_{r \notin \langle \xi \rangle} |a_r|^4$. If $b > 0$, the stronger finite inequality is

$$u^8 \geq 2S^2 + 12\rho^4 D + 16\rho^6 b^2 + 4 \left(2\rho^3 b |\cos \phi| - D_{\text{out}} \right)_+^2. \quad (\text{III.52.3})$$

Proof. In the four-frequency zero-sum representation of $\mathcal{E}(\lambda)$, retain, in addition to the three opposite-pair families, all permutations of $(\xi, \xi, \xi, 2\xi)$ and of its negative. There are eight such ordered tuples, each with weight $\rho^6 b^2$. They sum to zero because $5\xi = 0$, and none belongs to a retained opposite-pair family. Thus

$$u^8 \geq 5S^2 - 6M + 16\rho^6 b^2.$$

Conjugate pairing in the tail gives $M \leq 2\rho^8 + D^2/2$. Expanding $S = 2\rho^4 + D$ now yields

$$5S^2 - 6M \geq 8\rho^8 + 20\rho^4 D + 2D^2 = 2S^2 + 12\rho^4 D.$$

For the stronger statement retain the four off-axis slots $(s, t) = (\pm 2\xi, \pm 2\xi)$ in (III.50.3). They are distinct, and all have the same T value. The contribution from bases in $\langle \xi \rangle$ is

$$T(2\xi, 2\xi) = a_\xi^3 a_{2\xi} + \overline{a_\xi^3 a_{2\xi}} + E = 2\rho^3 b \cos \phi + E.$$

For a base outside $\langle \xi \rangle$, all four shifted indices remain outside it. Hölder's inequality therefore gives $|E| \leq D_{\text{out}}$. The four retained squares contribute at least the final term in (III.52.3). $\square$

The extra term $16\rho^6 b^2$ is the decisive order-five resonance. Without such a quadratic cost in the second harmonic, balancing a quartic tail cost against the quintic gain can produce fractional powers, as in the earlier $u^{16/3}$ result. Here the balance is quadratic and forces the smaller second harmonic $b \asymp u^2/\delta$.

Proposition III.52.2 (Explicit finite Fourier-profile bound). *For every $\delta \geq 0$ and every nonzero real centered f on an exponent-five group,*

$$\begin{aligned} Q_3(\delta + f) &\leq \delta^8 + 12\delta^4 S + 16\delta^3 \rho^4 b \cos \theta + 8\delta^3 \rho D \\ &\quad + 42\sqrt{3} \delta^2 u^4 D^{1/2} + 8\sqrt{3\sqrt{3}/2} \delta u^6 D^{1/4} + u^8. \end{aligned} \quad (\text{III.52.4})$$

If $b = 0$, its cosine term is interpreted as zero. This bound and the constraints (III.52.2)–(III.52.3) are finite inequalities, with no asymptotic qualification.

Proof. The two selected summands $r = \pm \xi$ in (III.51.2) give $2\rho^4 b \cos \theta$ exactly. Every other summand has absolute value at most $\rho |a_r|^4$, because ρ is the largest Fourier amplitude. Hence

$$B_5(f) = 2\rho^4 b \cos \theta + E_5, \quad |E_5| \leq \rho D. \quad (\text{III.52.5})$$

Use this identity and Theorem III.51.2 in (III.51.1). $\square$

III.53 Sharp uniform optimization and the phase envelope

Source 21, delivered Section 5.

We now prove Theorem F.26.1. The proof first locates the Fourier mass and then absorbs the remaining higher-degree terms. This order is necessary: dropping the negative tail term too early would lose the eighth-order remainder.

Lemma III.53.1 (Uniform localization). *Let $\|g\|_{U^3} = 1$, let $t > 0$, and put*

$$S_0 = 2^{-1/2}, \quad \eta = S_0 - S, \quad F = Q_3(1 + tg) - 1 - 6\sqrt{2}t^4,$$

where the Fourier statistics refer to g . If $F > 0$ and t is sufficiently small, then, with absolute implied constants,

$$\eta = O(t^2), \quad D = O(t^2), \quad b = O(t). \quad (\text{III.53.1})$$

Proof. First $S \leq S_0$ by Lemma III.50.2. Cube Cauchy–Schwarz and (III.51.1) give

$$F = -12\eta t^4 + 8B_5(g)t^5 + R, \quad |R| \leq 28t^6 + 8t^7 + t^8 \leq 37t^6 \quad (t \leq 1).$$

Also $|B_5(g)| \leq \rho S \leq 1$. Thus $F > 0$ implies $\eta = O(t)$, and S is bounded below by a positive constant for small t . Equation (III.50.7) yields

$$3DS \leq 1 - 2S^2 = 2\eta(S_0 + S),$$

so $D = O(\eta)$. It follows that $\rho^4 = (S_0 - \eta - D)/2$ is bounded below. The resonance penalty in (III.52.2) now gives $b^2 = O(\eta)$.

Equation (III.52.5) gives $|B_5(g)| = O(\sqrt{\eta} + \eta)$. Substituting in the preceding expression for F gives

$$\eta \leq At\sqrt{\eta} + Bt\eta + Ct^2$$

with absolute constants. For small t , absorb $Bt\eta$. The resulting quadratic inequality in $\sqrt{\eta}$ implies $\sqrt{\eta} = O(t)$, proving all three claims. $\square$

Proof of Theorem F.26.1, upper bound. If $u = 0$, the norm property gives $f = 0$. Otherwise set $g = f/u$ and $t = u/\delta$. Then $Q_3(\delta + f) = \delta^8 Q_3(1 + tg)$. If $F \leq 0$, the asserted upper bound is immediate, so assume $F > 0$. All the following Fourier variables refer to g .

By localization,

$$\rho^4 = 2^{-3/2} + O(t^2), \quad \rho^6 = 2^{-9/4} + O(t^2).$$

Expand (III.52.2), using $1 - 2S^2 = 2\sqrt{2}\eta - 2\eta^2$, to obtain

$$\eta \geq \frac{3}{2}D + 2^{1/4}b^2 - O(t^4). \quad (\text{III.53.2})$$

The constants are absolute. Use $B_5(g) \leq S_0 b + D$ and (III.51.6) to get

$$\begin{aligned} F \leq & -18Dt^4 - 12 \cdot 2^{1/4}b^2t^4 + 4\sqrt{2}bt^5 + 8Dt^5 \\ & + K_6t^6D^{1/2} + K_7t^7D^{1/4} + O(t^8), \end{aligned} \quad (\text{III.53.3})$$

where $K_6 = 42\sqrt{3}$ and $K_7 = 8\sqrt{3\sqrt{3}/2}$.

For sufficiently small t , $-18D + 8Dt \leq -17D$. Set $D = t^4y$. All the terms involving D in (III.53.3) are then bounded above by

$$t^8(-17y + K_6\sqrt{y} + K_7y^{1/4}) = O(t^8),$$

because the continuous expression in parentheses tends to $-\infty$ as $y \rightarrow \infty$. Finally,

$$\sup_{b \geq 0} \{-12 \cdot 2^{1/4} b^2 t^4 + 4\sqrt{2} b t^5\} = C_5 t^6, \quad b = \beta_5 t. \quad (\text{III.53.4})$$

Rescaling by δ^8 proves (F.26.2). The error is $O(\delta^8 t^8) = O(u^8)$ with an absolute constant, not merely an unspecified density-dependent error. $\square$

Theorem III.53.2 (Sharp phase-dependent coefficient). *Under the hypotheses of Theorem F.26.1, choose a largest conjugate pair. If $b > 0$ and the phases are as in (III.52.1), then, with absolute $c, C > 0$,*

$$Q_3(\delta + f) \leq \delta^8 + 6\sqrt{2} \delta^4 u^4 + C_5 \frac{(\cos \theta)_+^2}{1 + \cos^2 \phi} \delta^2 u^6 + C u^8 \quad (u \leq c\delta). \quad (\text{III.53.5})$$

For $b = 0$, the same assertion holds with the u^6 coefficient zero. For each prescribed pair of phases, the coefficient in (III.53.5) is attained asymptotically by weights on $\mathbb{Z}/5\mathbb{Z}$.

Proof. As before, the case $F \leq 0$ is immediate. In the case $F > 0$, localization applies. The elementary inequality $(x - y)_+^2 \geq x^2 - 2xy$ for $x, y \geq 0$ turns (III.52.3) into

$$\eta \geq \frac{3}{2} D + 2^{1/4} (1 + \cos^2 \phi) b^2 - O(t^4 + bD).$$

The selected quintic term is $2\rho^4 b \cos \theta$. Replacing $2\rho^4$ by S_0 here introduces an absolute error $O((\eta + D)b) = O(t^3)$, regardless of the sign of the cosine; after multiplication by t^5 this is $O(t^8)$. The extra $O(bD)t^4$ is absorbed against the negative tail term because $b = O(t)$. The other tail terms are treated exactly as in (III.53.3).

The remaining one-variable optimization is

$$\sup_{b \geq 0} \{-12 \cdot 2^{1/4} (1 + \cos^2 \phi) b^2 t^4 + 4\sqrt{2} b \cos \theta t^5\} = C_5 \frac{(\cos \theta)_+^2}{1 + \cos^2 \phi} t^6.$$

The maximizing scale is

$$b = \beta_5 t \frac{(\cos \theta)_+}{1 + \cos^2 \phi}. \quad (\text{III.53.6})$$

For fixed phases with $\cos \theta > 0$, the exact five-point formulas in Section III.55, with unnormalized second amplitude $\beta_5 \delta^{-1} u^2 (\cos \theta)_+ / (1 + \cos^2 \phi)$, give a matching coefficient and an $O_\delta(u^8)$ error. The first amplitude is chosen to make the centered norm exactly u . If $\cos \theta \leq 0$, take a positive second amplitude $o(\delta^{-1} u^2)$ with the specified phase. The coefficient is then zero; for example $b = \delta^{-3} u^4$ gives an $O_\delta(u^8)$ error. The case $b = 0$ follows directly from the same absorption argument with no selected quintic term. $\square$

Remark III.53.3 (Real Fourier coefficients lose a factor of two). For a positive real dominant coefficient and a positive real second coefficient, $\theta = \phi = 0$. The phase factor in (III.53.5) is $1/2$. The full coefficient requires $\theta = 0$ and $\cos \phi = 0$, which can be achieved by $a_\xi = \rho e^{i\pi/10}$ and $a_{2\xi} = b e^{i\pi/5}$. The extremizing physical-space function is still real: the negative Fourier coefficients are the complex conjugates of the positive ones.




Figure III.53.1: The exact phase factor in Theorem III.53.2. The marked point $\theta = 0, \phi = \pi/2$ gives the full sharp coefficient; positive real coefficients give the value $1/2$ at $\theta = \phi = 0$. The factor vanishes when $\cos \theta \leq 0$.

III.54 Near equality and the scale of rigidity

Source 21, delivered Section 6.

Theorem III.54.1 (Second-term rigidity). *Fix $\delta > 0$. Suppose G_j are nontrivial exponent-five groups, f_j are real and centered, $0 < u_j = \|f_j\|_{U^3} \rightarrow 0$, and*

$$Q_3(\delta + f_j) = \delta^8 + 6\sqrt{2}\delta^4 u_j^4 + C_5 \delta^2 u_j^6 + o_\delta(u_j^6). \quad (\text{III.54.1})$$

For any choice of a largest conjugate pair $\{\pm \xi_j\}$, write $\hat{f}_j(\xi_j) = \rho_j e^{i\alpha_j}$ and $\hat{f}_j(2\xi_j) = b_j e^{i\beta_j}$. Then

$$\begin{aligned} b_j &\sim \beta_5 \delta^{-1} u_j^2, & e^{i(2\alpha_j - \beta_j)} &\longrightarrow 1, \\ e^{10i\alpha_j} &\longrightarrow -1, & \sum_{r \notin \{\pm \xi_j\}} |\hat{f}_j(r)|^4 &= o_\delta(u_j^6), \\ \rho_j^4 &= \frac{u_j^4}{2\sqrt{2}} - \frac{C_5}{24} \delta^{-2} u_j^6 + o_\delta(u_j^6). \end{aligned} \quad (\text{III.54.2})$$

Consequently the remainder after removing the dominant pair is $o_\delta(u_j^{3/2})$ in U^2 . This conclusion does not assert an $o(u_j^2)$ remainder or an L^2 approximation.

Proof. Normalize $g_j = f_j/u_j$, $t_j = u_j/\delta$. All variables in this paragraph refer to g_j . The localization lemma gives $D = O(t^2)$ and $b = O(t)$. Dividing the upper argument by t^6 , and retaining the phase and tail losses, gives

$$\frac{F}{t^6} \leq -18 \frac{D}{t^2} - 12 \cdot 2^{1/4} \left(\frac{b}{t}\right)^2 + 4\sqrt{2} \frac{b}{t} \cos \theta + o(1).$$

Here the sextic and septic errors are respectively $O(t)$ and $O(t^{3/2})$ after division, so both tend to zero. The quadratic has unique maximum C_5 at $b/t = \beta_5$, $\cos \theta = 1$. Thus $D/t^2 \rightarrow 0$, $b/t \rightarrow \beta_5$, and $e^{i\theta} \rightarrow 1$.

The phase envelope then forces $\cos \phi \rightarrow 0$, since its factor must tend to one. As $\theta + \phi = 5\alpha$, it follows that $e^{10i\alpha} \rightarrow -1$.

The exact expansion gives

$$C_5 = -12 \lim \frac{\eta}{t^2} + 4\sqrt{2}\beta_5.$$

Since $4\sqrt{2}\beta_5 = 2C_5$, this yields $\eta/t^2 \rightarrow C_5/12$. Finally, $\rho^4 = (S_0 - \eta - D)/2$ gives the last line of (III.54.2). Restore the powers of u_j and δ . The Fourier identity for U^2 identifies the fourth power of the remainder norm with the displayed tail mass. $\square$

[write] **6 October 2026; sources 21–27.** *The same conclusions in source 23* (checked at the write). Source 23's Theorem III.58.3 writes $a_\xi = \rho e^{i\theta}$, $a_{2\xi} = b e^{i\psi}$ (its θ, ψ are this theorem's α, β) and $\kappa = 2^{1/4}/6 = \beta_5$. Its $e^{i(\psi-2\theta)} \rightarrow 1$ is the second limit here; its $\cos(5\theta) \rightarrow 0$ is equivalent to $e^{10i\alpha} \rightarrow -1$; and its last line $\rho^4 = 2^{-3/2}u^4 - 2^{-3/4}\kappa^2\delta^{-2}u^6 + o(u^6)$ is the last line here, since $2^{-3/4}\kappa^2 = 2^{-1/4}/36 = C_5/24$.

III.54.1 Why the remainder scale cannot be silently strengthened

The forced second harmonic has amplitude of order u^2 , while the tail conclusion in Theorem III.54.1 is only $o(u^{3/2})$ in U^2 . These statements are consistent. A further independent harmonic of amplitude u^a , with $3/2 < a < 2$, can have fourth-power mass $o(u^6)$ while being larger than u^2 . Such a harmonic is invisible to extremality through order u^6 . To obtain sharper Fourier approximation one must control the objective at the next order as well.

The indicator examples in Section III.56 give a stronger obstruction to L^2 rigidity. For them the centered L^2 norm tends to $\sqrt{\delta(1-\delta)}$, even though every fixed finite collection of Fourier coefficients tends to zero. Thus the appropriate topology for the stated rigidity conclusion is essential.

III.55 The exact next coefficient on the five-point group

Source 21, delivered Section 7.

Let $G = \mathbb{Z}/5\mathbb{Z}$, and use normalized averages. For real mean-zero $f : G \rightarrow \mathbb{R}$, put $u = \|f\|_{U^3}$ and

$$Q_3(W) = \mathbb{E}_{x, h_1, h_2, h_3} \prod_{\omega \in \{0,1\}^3} W(x + \omega \cdot h).$$

For fixed $0 < \delta < 1$, let $M_{\delta,5}(u)$ be the maximum of $Q_3(W)$ over $W : G \rightarrow [0,1]$ with mean δ and $\|W - \delta\|_{U^3} = u$. The same conclusions hold with the boundedness restriction omitted and with any fixed $\delta > 0$.

Theorem III.55.1 (Third extremal term on $\mathbb{Z}/5\mathbb{Z}$). *As $u \downarrow 0$,*

$$M_{\delta,5}(u) = \delta^8 + 6\sqrt{2}\delta^4u^4 + \frac{2^{3/4}}{3}\delta^2u^6 + \frac{20}{9}u^8 + O_\delta(u^{10}).$$

The assertion concerns this fixed five-point group. It does not identify the coefficient of u^8 when the dimension of an exponent-five group is allowed to vary.

[write] **6 October 2026; sources 21–27.** *Credit and scope.* Source 23's Appendix A (Section III.59) computes the coefficient 20/9 for one five-point family, which bounds the coefficient of the maximum from below; the theorem here proves that 20/9 is the coefficient of the maximum on $\mathbb{Z}/5\mathbb{Z}$. The intake's independent

maximization on $\mathbb{Z}/5\mathbb{Z}$ (50-digit arithmetic) gave normalized eighth-order residuals 2.32462, 2.24655, 2.22823, 2.22372 at $t = 0.2, 0.1, 0.05, 0.025$, converging to $20/9 = 2.2222\dots$ with an $O(t^2)$ gap; that is evidence only. The universal coefficient, over groups of growing dimension, is open (item III.27).

III.55.1 An exact Fourier certificate with arbitrary phases

Write

$$\widehat{f}(1) = \rho e^{i\alpha}, \quad \widehat{f}(2) = b e^{i\beta}, \quad \widehat{f}(-r) = \overline{\widehat{f}(r)},$$

where $\rho, b \geq 0$, and set

$$\theta = 2\alpha - \beta, \quad \phi = 3\alpha + \beta, \quad \psi = \phi - \theta.$$

Interchanging the two conjugate frequency pairs, if necessary, makes $\rho \geq b$. For $a_r = \widehat{W}(r)$ define

$$T(s, t) = \sum_{r \in G} a_r \overline{a_{r+s}} \overline{a_{r+t}} a_{r+s+t}.$$

Fourier expansion and orthogonality give $Q_3(W) = \sum_{s,t \in G} |T(s, t)|^2$. Because W is real, substitution $r \mapsto -r - s - t$ shows that $T(s, t)$ is real; shifts of r and interchange of s, t show that its value is unchanged under independent sign changes of s, t and their interchange. The following six representatives therefore exhaust all 25 slots:

(s, t)	multiplicity	$T(s, t)$
$(0, 0)$	1	$\delta^4 + 2\rho^4 + 2b^4$
$(0, 1)$	4	$2\delta^2\rho^2 + 2\rho^2b^2 + b^4$
$(0, 2)$	4	$2\delta^2b^2 + 2\rho^2b^2 + \rho^4$
$(1, 1)$	4	$\delta^2\rho^2 + 2\delta\rho^2b \cos \theta + 2\rho b^3 \cos(2\theta - \phi)$
$(2, 2)$	4	$\delta^2b^2 + 2\delta\rho b^2 \cos \psi + 2\rho^3b \cos \phi$
$(1, 2)$	8	$\rho^2b^2 + 2\delta\rho^2b \cos \theta + 2\delta\rho b^2 \cos \psi.$

Each row is obtained by substituting the five possible values of r . In particular, the exact centered norm is

$$\begin{aligned} N(\rho, b, \theta, \phi) := u^8 &= 8\rho^8 + 16(1 + \cos^2 \phi)\rho^6b^2 + 48\rho^4b^4 \\ &\quad + 16(1 + \cos^2(2\theta - \phi))\rho^2b^6 + 8b^8. \end{aligned} \quad (\text{III.55.1})$$

Squaring the rows of the table gives the following complete formula:

$$\begin{aligned} Q_3(\delta + f) &= \delta^8 + 24\delta^4(\rho^4 + b^4) \\ &\quad + 16\delta^3(\rho^4b \cos \theta + \rho b^4 \cos \psi) \\ &\quad + 24\delta^2\left((3 + \cos 2\theta)\rho^4b^2 + (3 + \cos 2\psi)\rho^2b^4\right. \\ &\quad \quad \left.+ 2\rho^3b^3(\cos \phi + \cos(2\theta - \phi))\right) \\ &\quad + 16\delta\left(\rho^4b^3(3 \cos \theta + \cos(2\phi - \theta))\right. \\ &\quad \quad \left.+ \rho^3b^4(3 \cos \psi + \cos(3\theta - \phi))\right) + N(\rho, b, \theta, \phi). \end{aligned} \quad (\text{III.55.2})$$

Thus the amplitude penalty $16\rho^6b^2$ cannot be decreased by a different phase choice. It is minimized by $\cos \phi = 0$, which is compatible with maximizing the leading quintic contribution through $\cos \theta = 1$.

III.55.2 Localization of every maximizer

We first optimize without the interval restriction. By homogeneity it suffices to take $\delta = 1$ and write $t = u/\delta$. For a five-point function the norm sphere is compact, so a maximum exists. More explicitly, the nonnegative coefficients in (III.55.1) and $\rho \geq b$ give

$$8\rho^8 \leq t^8 \leq 128\rho^8.$$

Consequently $\rho \asymp t$, uniformly in the phases, and $\|f\|_\infty \leq 2(\rho + b) = O(t)$. After rescaling back to fixed $0 < \delta < 1$, the interval restriction $0 \leq \delta + f \leq 1$ is therefore inactive for all sufficiently small u .

Put $S = 2(\rho^4 + b^4)$. The exact norm gives

$$t^8 - 2S^2 \geq 16\rho^6 b^2,$$

and hence

$$6\sqrt{2}t^4 - 24(\rho^4 + b^4) = \frac{6(t^8 - 2S^2)}{t^4/\sqrt{2} + S} \geq ct^2b^2$$

for an absolute $c > 0$. Since $b \leq \rho = O(t)$, all the quintic, sextic, and septic terms in (III.55.2) are bounded in absolute value, respectively, by Ct^4b , Ct^4b^2 , and Ct^4b^3 . In particular, for small t ,

$$Q_3(1+f) - 1 - 6\sqrt{2}t^4 \leq -c't^2b^2 + Ct^4b + t^8.$$

The pure-pair example $b = 0$ gives the value $1 + 6\sqrt{2}t^4 + t^8$. A maximizing function must therefore have $b \leq C't^2$. This is the needed compactness statement at the secondary scale: write $b = kt^2$, with k bounded independently of t .

III.55.3 The leading phase optimization

Set $q = 2^{-3/4}$ and $\rho^2 = t^2 R$. After dividing the exact norm by t^8 , its defining equation is

$$\begin{aligned} 1 &= 8R^4 + 16(1 + \cos^2 \phi)R^3 k^2 t^2 + 48R^2 k^4 t^4 \\ &\quad + 16(1 + \cos^2(2\theta - \phi))Rk^6 t^6 + 8k^8 t^8. \end{aligned}$$

At $t = 0$, its positive solution is $R = q$. The derivative with respect to R is $32q^3 > 0$, so the implicit function theorem gives a real analytic solution $R(t, k, \theta, \phi)$, uniformly on each bounded set of k . In fact it is analytic in t^2 .

Subtract the constant and quartic terms from the objective and divide by t^6 . The apparent singularity at $t = 0$ is removable. The resulting real analytic function H has the uniform expansion

$$H(t, k, \theta, \phi) = H_0(k, \theta, \phi) + t^2 H_2(k, \theta, \phi) + t^3 H_3(k, \theta, \phi) + O(t^4),$$

where

$$H_0 = -24q(1 + \cos^2 \phi)k^2 + 16q^2 k \cos \theta.$$

For $k \geq 0$, the maximum is

$$\max H_0 = \frac{8q^3}{3} = \frac{2^{3/4}}{3},$$

and equality holds precisely at

$$k_0 = q/3, \quad \theta = 0 \pmod{2\pi}, \quad \phi = \pi/2 \pmod{\pi}.$$

Indeed, $H_0 \leq -24qk^2 + 16q^2k$, with equality forcing those phases when $k > 0$, and the remaining quadratic has its unique maximum at k_0 . Compactness and uniform convergence imply that every maximizer for sufficiently small t lies near one of these two phase classes. At either class the Hessian, in coordinates (k, θ, ϕ) , is

$$\text{diag} \left(-48q, -\frac{16q^3}{3}, -\frac{16q^3}{3} \right).$$

It is negative definite. Applying the implicit function theorem to $\nabla H = 0$ produces one analytic local maximum near each class; these exhaust the global maxima. Complex conjugation interchanges the two classes and preserves the objective, so their values agree. The map $(\alpha, \beta) \mapsto (\theta, \phi)$ has determinant 5; its finite collection of lifts merely produces the translates of these same functions on the five-point group.

III.55.4 Calculation of the coefficient 20/9

At $\theta = 0$, $\phi = \pi/2$, equation (III.55.1) becomes

$$t^8 = 8\rho^8 + 16\rho^6b^2 + 48\rho^4b^4 + 16\rho^2b^6 + 8b^8.$$

Let $x = t^2$, $b = kx$, and write $\rho^2 = qxX(z)$, where $z = k^2x/q$. Then

$$X^4 + 2zX^3 + 6z^2X^2 + 2z^3X + z^4 = 1, \quad X(0) = 1.$$

Successive coefficient comparison gives

$$X(z) = 1 - \frac{1}{2}z - \frac{9}{8}z^2 + \frac{3}{4}z^3 + O(z^4),$$

and therefore

$$\rho^4 = q^2t^4 - qk^2t^6 - 2k^4t^8 + O(t^{10}).$$

The exact objective at these phases simplifies to

$$Q_3(1 + f) = 1 + 24(\rho^4 + b^4) + 16\rho^4b + 96\rho^4b^2 + 48\rho^2b^4 + 32\rho^4b^3 + t^8.$$

It follows that

$$\begin{aligned} H_0(k, 0, \pi/2) &= -24qk^2 + 16q^2k, \\ H_2(k, 0, \pi/2) &= 1 - 24k^4 - 16qk^3 + 96q^2k^2. \end{aligned}$$

At $k_0 = q/3$, using $q^4 = 1/8$, we obtain

$$H_2(k_0, 0, \pi/2) = 1 - \frac{1}{27} - \frac{2}{27} + \frac{4}{3} = \frac{20}{9}.$$

It remains to justify the claimed $O(t^{10})$ remainder, rather than merely $o(t^8)$. The analytic maximizing coordinates differ from their $t = 0$ limits by $O(t^2)$ because H has no term of order t and the Hessian of H_0 is invertible. Substitution changes H_0 by $O(t^4)$ and changes t^2H_2 by $O(t^4)$. The only possible contribution of order t^3 in H is

$$H_3(k, \theta, \phi) = 16\sqrt{q}k^4 \cos(\phi - \theta) + 48q^{3/2}k^3(\cos(2\theta - \phi) + \cos \phi).$$

It vanishes at both limiting phase classes. Thus its contribution at the moving maximizers is $O(t^5)$, and in particular $O(t^4)$. Consequently

$$\max H = \frac{2^{3/4}}{3} + \frac{20}{9}t^2 + O(t^4).$$

Multiplying by t^6 , restoring the constant and quartic terms, and rescaling by δ^8 proves Theorem III.55.1.

What the optimizing phases do. The limiting phase choice is sufficient to calculate the terms through u^8 , but it need not remain exactly optimal for positive u . The t^3H_3 perturbation can move the phases. The theorem therefore uses the analytic optimizing branch, and does not assume exact phase stationarity of the displayed elementary model away from $t = 0$.

Reproducible certificate. The accompanying Python program `verify_order5.py` independently enumerates the 5^4 solutions of the Fourier cube constraints, records their monomials, and checks the six-slot formula against direct averages over 5^4 physical cubes. It also checks the rational series coefficients and the evaluation $20/9$. These finite computations verify identities; the maximization and the asymptotic remainder follow from the argument above.

III.55.5 Sharpness on every exponent-five group

Corollary III.55.2 (Pullback of the extremizing model). *Let G be any fixed nontrivial finite abelian group of exponent five and $0 < \delta < 1$. There are $[0, 1]$ weights W_u on G , with mean δ and centered norm exactly u , for all sufficiently small $u > 0$, such that*

$$Q_3(W_u) = \delta^8 + 6\sqrt{2}\delta^4u^4 + C_5\delta^2u^6 + \frac{20}{9}u^8 + O_\delta(u^{10}). \quad (\text{III.55.3})$$

Consequently the uniform coefficient in Theorem F.26.1 is optimal on each such fixed group.

Proof. Choose a nonzero linear functional $\pi : G \rightarrow \mathbb{F}_5$. It is surjective, and the image of a uniform element of G is uniform on $\mathbb{F}_5$. For any weight V on $\mathbb{F}_5$, the pullback $V \circ \pi$ has the same mean, centered U^3 norm, and normalized cube count as V , since the four cube parameters push forward independently and uniformly. Pull back the five-point maximizers from Theorem III.55.1. $\square$

[write] 6 October 2026; sources 34–39. *Dimension-uniform* (source 38). On every group of exponent five the maximum reduces exactly to $\mathbb{Z}/5\mathbb{Z}$ for $u \leq 10^{-5}\delta$, so $20/9$ is the universal coefficient of u^8 (source 38's Theorem III.84.1); source 38 also gives the u^{10} and u^{12} coefficients (Theorem III.85.1).

For the coefficients through u^8 , an elementary explicit model already suffices. Set

$$b = \beta_5\delta^{-1}u^2, \quad a_1 = \rho e^{i\pi/10}, \quad a_2 = be^{i\pi/5}, \quad (\text{III.55.4})$$

with conjugate coefficients at $-1, -2$, and choose the unique positive ρ satisfying

$$8\rho^8 + 16\rho^6b^2 + 48\rho^4b^4 + 16\rho^2b^6 + 8b^8 = u^8. \quad (\text{III.55.5})$$

Existence and uniqueness for small u follow from strict monotonicity in ρ and from $8b^8 < u^8$. In physical space this is

$$W_u(x) = \delta + 2\rho \cos\left(\frac{2\pi x}{5} + \frac{\pi}{10}\right) + 2b \cos\left(\frac{4\pi x}{5} + \frac{\pi}{5}\right).$$




Figure III.55.1: The exact elementary model at $\delta = 0.4$, evaluated with high-precision arithmetic. The horizontal axis is $t = u/\delta$. Left: the normalized sixth-order correction tends to C_5 . Right: after subtracting that correction as well, the normalized eighth-order residual tends to $20/9$. The plot evaluates a proven formula; it is not used as evidence for the general upper bound.

The coefficients tend to zero, so $W_u \in [0, 1]$ when u is small. The fixed choice of the limiting amplitude and phases has the same $20/9$ coefficient, although its terms beyond u^8 need not be globally optimal. Figure III.55.1 illustrates the approach to the two sharp coefficients using the exact formulas.

III.56 Indicator sharpness and its limitations

Source 21, delivered Section 8.

Weights are convenient because the Fourier parameters vary continuously. The sharp coefficients also have consequences for actual sets. A fixed five-point group has only finitely many sets, so a nontrivial small-norm limit at a prescribed positive density requires the group to grow.

Theorem III.56.1 (Sharp indicator families). *For every $0 < \delta < 1$, there are dimensions $n_j \rightarrow \infty$ and sets $A_j \subset \mathbb{F}_5^{n_j}$ with densities $\delta_j \rightarrow \delta$ and centered norms $v_j = \|1_{A_j} - \delta_j\|_{U^3} \rightarrow 0$ such that*

$$Q_3(1_{A_j}) = \delta_j^8 + 6\sqrt{2}\delta_j^4 v_j^4 + C_5 \delta_j^2 v_j^6 + \frac{20}{9} v_j^8 + o(v_j^8). \quad (\text{III.56.1})$$

Thus C_5 remains optimal for indicator sets in growing exponent-five groups, and $20/9$ is a lower bound for any universal eighth-order extremal coefficient in that setting.

Proof. Choose $u_j \downarrow 0$ and lift the elementary model (III.55.4) to $G_j = \mathbb{F}_5^{n_j}$ by its first coordinate. Denote the lifted weight by W_j . Its mean is δ , its centered norm is u_j , and its cube count has expansion (III.55.3) through u_j^8 , independently of n_j .

Include each point $x \in G_j$ independently with probability $W_j(x)$. Let X_x be its indicator. We approximate three statistics: the density, $Q_3(X)$, and $Q_3(X - \delta)$. If the eight vertices of a cube are distinct, independence makes the expectation of its product equal to the corresponding weighted product. For any two distinct vertices, the equation that they coincide is a nontrivial linear equation

in the sidelengths and holds with probability $|G_j|^{-1}$. Hence at most $28/|G_j|$ of all normalized cubes have a repeated vertex. Since all relevant factors have absolute value at most one, the expectation error in either cube statistic is at most $56/|G_j|$.

Changing one inclusion variable affects at most $8/|G_j|$ of either normalized cube polynomial, by a union bound over the vertex at which that point occurs. A bounded-difference variance estimate, or the elementary Efron–Stein inequality, gives variance $O(|G_j|^{-1})$. The density has variance at most $1/(4|G_j|)$. Chebyshev’s inequality and a union bound therefore show that, for sufficiently large n_j , some realization simultaneously approximates all three target statistics to an arbitrarily prescribed positive accuracy.

Choose the accuracy to be $o(u_j^{12})$, for example a sufficiently small multiple of u_j^{14} . Recentring X from δ to its actual density δ_j changes a cube product by at most $8|\delta_j - \delta|$, by telescoping. Thus

$$v_j^8 = u_j^8 + o(u_j^{12}), \quad \delta_j - \delta = o(u_j^{12}), \quad Q_3(1_{A_j}) - Q_3(W_j) = o(u_j^{12}).$$

It follows that $v_j/u_j \rightarrow 1$ and $v_j^4 - u_j^4 = o(u_j^8)$; the corresponding changes in the sixth- and eighth-order terms are also $o(u_j^8)$. Substitution proves (III.56.1). The dimensions can be chosen strictly increasing at each step. $\square$

[write] **6 October 2026; sources 21–27.** *Credit.* Source 23’s Proposition III.58.4 realizes the coefficients through order u^6 by indicator sets in the same way; the theorem here goes through u^8 . Both are lower examples, not upper bounds.

Corollary III.56.2 (A strengthened set version of Lemma 3.10). *Let $A \subset G$, where G is a nontrivial finite abelian group of exponent five, let $N = |G|$ and $\delta = |A|/N > 0$, and suppose*

$$\|1_A - \delta\|_{U^3}^8 \leq \alpha.$$

For absolute $c, C > 0$, if $\alpha^{1/8} \leq c\delta$, then

$$\#\{\text{parameterized three-cubes in } A\} \leq N^4 \left(\delta^8 + 6\sqrt{2}\delta^4\alpha^{1/2} + C_5\delta^2\alpha^{3/4} + C\alpha \right). \quad (\text{III.56.2})$$

Proof. Apply Theorem F.26.1 at the actual centered norm and then increase that norm to its budget $\alpha^{1/8}$. All displayed powers have positive coefficients, so this substitution preserves the upper bound. Multiply the normalized count by N^4 . $\square$

In this setting the triangle-inequality estimate corresponding to Gowers’s Lemma 3.10 is $N^4(\delta + \alpha^{1/8})^8$. The first error in (III.56.2) is order $\alpha^{1/2}$, and its next sharp coefficient occurs at $\alpha^{3/4}$. The characteristic-five restriction is essential to this particular secondary statement. It is not a replacement for the prime-cyclic bound at arbitrary primes.

III.56.1 Failure of stronger L^2 rigidity

For the indicator examples,

$$\|1_{A_j} - \delta_j\|_2^2 = \delta_j(1 - \delta_j) \longrightarrow \delta(1 - \delta) > 0.$$

On the other hand each Fourier coefficient is $O(v_j)$, by $\|1_{A_j} - \delta_j\|_{U^2}^4 \leq v_j^4/\sqrt{2}$. If P_j projects onto at most a fixed number m of frequencies, Parseval gives $\|P_j(1_{A_j} - \delta_j)\|_2^2 = O(mv_j^2) \rightarrow 0$. The orthogonal remainder therefore has L^2 norm tending to $\sqrt{\delta(1 - \delta)}$. Even the strong asymptotic equality in (III.56.1) does not force a small L^2 remainder after removing finitely many harmonics.

III.57 A quantitative projection consequence of cube Cauchy–Schwarz

Source 21, delivered Section 9.

The proof of the sharp coefficient used the stronger explicit sixth- and seventh-degree formulas. For related problems where such formulas are unavailable, a general norm argument still provides a uniform way to pass from near equality to approximation on a quotient. The following is an elementary consequence of classical norm geometry; its inclusion is methodological, not a priority claim.

Proposition III.57.1 (Loss under subgroup averaging controls the remainder). *Let G be any finite abelian group, $d \geq 2$, $p = 2^d$, and*

$$P_H f(x) = \mathbb{E}_{h \in H} f(x + h)$$

for a subgroup $H \leq G$. Then

$$\|f - P_H f\|_{U^d} \leq 2 \left(\|f\|_{U^d}^p - \|P_H f\|_{U^d}^p \right)^{1/p}. \quad (\text{III.57.1})$$

In particular, equality of the two norms holds if and only if $f = P_H f$.

Proof. Cube expansion and Gowers–Cauchy–Schwarz yield, with $a = \|f\|_{U^d}$ and $b = \|g\|_{U^d}$,

$$\|f + g\|_{U^d}^p + \|f - g\|_{U^d}^p \leq 2 \sum_{j \text{ even}} \binom{p}{j} a^{p-j} b^j.$$

Use weighted AM–GM on each monomial and the identities

$$\sum_{j \text{ even}} j \binom{p}{j} = \sum_{j \text{ even}} (p - j) \binom{p}{j} = p 2^{p-2}$$

to obtain Clarkson’s inequality

$$\left\| \frac{f + g}{2} \right\|_{U^d}^p + \left\| \frac{f - g}{2} \right\|_{U^d}^p \leq \frac{\|f\|_{U^d}^p + \|g\|_{U^d}^p}{2}. \quad (\text{III.57.2})$$

Subgroup averaging is contractive by translation invariance and the triangle inequality. Apply (III.57.2) to f and its translate $T_h f$ for $h \in H$. Since $P_H((f + T_h f)/2) = P_H f$, the first summand on the left is at least $\|P_H f\|_{U^d}^p$. Hence

$$\|f - T_h f\|_{U^d} \leq 2 \left(\|f\|_{U^d}^p - \|P_H f\|_{U^d}^p \right)^{1/p}.$$

Average $f - T_h f$ over $h \in H$ and use the triangle inequality. The final equivalence follows because U^d is a norm for $d \geq 2$. $\square$

For $d = 3$ and $H = \ker \xi$, P_H retains exactly the frequencies in $\langle \xi \rangle$. Thus near saturation of its norm implies approximation by a five-point quotient in U^3 , uniformly over the ambient dimension. The proposition alone does not establish that a global maximizer of $Q_3(\delta + f)$ lies on that quotient: projection also changes the centered norm, and restoring the norm by rescaling changes the different-degree terms unequally. This is precisely an obstacle in the higher-dimensional eighth-order question.

III.58 Sharp secondary three-cube asymptotics in characteristic five

Source 23, delivered Section 7.

The earlier project report *Sharp Extremal Refinements in Gowers's Szemerédi Framework* proved that exponent five changes the first secondary power from $u^{16/3}$ to u^6 [112]. Its explicit example had coefficient $2^{3/4}/3$, but optimality of this coefficient was left open. We determine that coefficient, uniformly in the dimension, and obtain an $O(u^8)$ remainder. These are local extremal statements about cube counts, with no claim of an improved global quantitative Szemerédi theorem.

For a real function W on a finite abelian group G , write

$$Q(W) = \mathbb{E}_{x, h_1, h_2, h_3 \in G} \prod_{v \in \{0,1\}^3} W(x + v \cdot h) = \|W\|_{U^3}^8.$$

All averages and Fourier transforms are normalized; sums on the dual are unnormalized. Repeated vertices and zero directions are included.

For comparison with Gowers's Lemma 3.10, his degree-two uniformity parameter is $\alpha = u^8$. His triangle-inequality bound has normalized form $(\delta + u)^8$ [1]. The sharper leading term $6\sqrt{2}\delta^4 u^4$ and the order-five lower example were established in the earlier project reports [108, 112]. The result below determines the previously unresolved secondary constant.

Theorem III.58.1 (The exact characteristic-five coefficient). *There are absolute constants $c, C > 0$ such that the following holds for every nontrivial finite abelian group of exponent five. If $\delta > 0$, $f : G \rightarrow \mathbb{R}$ has mean zero, and $u = \|f\|_{U^3} \leq c\delta$, then*

$$Q(\delta + f) \leq \delta^8 + 6\sqrt{2}\delta^4 u^4 + \frac{2^{3/4}}{3} \delta^2 u^6 + Cu^8. \quad (\text{F1})$$

No pointwise bound on f is required, and c, C are independent of the group and its dimension. For every fixed $0 < \delta < 1$, on every fixed nontrivial group G of exponent five, the supremum over weights $W : G \rightarrow [0, 1]$ satisfying $\mathbb{E}W = \delta$ and $\|W - \delta\|_{U^3} \leq u$ obeys

$$M_{\delta, G}(u) = \delta^8 + 6\sqrt{2}\delta^4 u^4 + \frac{2^{3/4}}{3} \delta^2 u^6 + O(u^8) \quad (u \downarrow 0). \quad (\text{F2})$$

The same expansion holds when the supremum also ranges over all such groups. The range of sufficiently small u in the lower construction may depend on δ , because the weights must lie in $[0, 1]$.

[write] **6 October 2026; sources 21–27.** *Second route.* This is source 21's Theorem F.26.1, proved independently by the same mechanism (a positive cubic statistic for the incomplete cubes, the quadratic cost of the second harmonic); source 21 has sharper intermediate constants, the phase envelope and the next coefficient on $\mathbb{Z}/5\mathbb{Z}$ (Section III.49). It answers source 14's Question 1 (item III.23).

Corollary III.58.2 (Ordered cube counts in uniform subsets). *There are absolute $c, C > 0$ such that, if $A \subset G = \mathbb{F}_5^n$ has density $\delta > 0$ and $\|1_A - \delta\|_{U^3}^8 \leq \alpha$ with $\alpha^{1/8} \leq c\delta$, its number of ordered three-cubes is at most*

$$|G|^4 \left(\delta^8 + 6\sqrt{2}\delta^4 \alpha^{1/2} + \frac{2^{3/4}}{3} \delta^2 \alpha^{3/4} + C\alpha \right).$$

Proof. Apply Theorem III.58.1 at the actual centered norm and use monotonicity of the displayed nonnegative terms. Normalization of the four independent cube parameters supplies the factor $|G|^4$. $\square$

III.58.1 The exact terms of degree six and a useful degree-seven bound

Let $a_r = \widehat{f}(r)$, so $a_0 = 0$ and $a_{-r} = \overline{a_r}$. Put

$$\lambda_r = |a_r|^2, \quad S = \sum_r \lambda_r^2 = \|f\|_{U^2}^4, \quad T(s, t) = \sum_r a_r \overline{a_{r+s}} \overline{a_{r+t}} a_{r+s+t}.$$

Fourier inversion and orthogonality give

$$u^8 = \sum_{s,t} |T(s, t)|^2. \quad (\text{F3})$$

One can see this identity directly by expanding all eight physical-space factors into characters. The surviving frequencies satisfy the four equations saying that the coefficient of x, h_1, h_2, h_3 is zero. After multiplying the frequency at vertex v by $(-1)^{|v|}$, their solutions are exactly affine cubes $r + sv_1 + tv_2 + zv_3$ in the dual group. This follows either by solving the four equations or by Boolean finite differences. Since $a_{-r} = \overline{a_r}$, the surviving product is the usual conjugated dual cube product. Summing over the last direction z pairs two translated square products, giving $\sum_{s,t} |T(s, t)|^2$ with no missing normalization factor. Define the following homogeneous forms:

$$\begin{aligned} B_5 &= \sum_r a_{2r} a_{-r}^3 a_r, \\ A_6 &= \sum_{s,t} \lambda_s \lambda_t \lambda_{s+t}, \\ B_6 &= \sum_{s,t} a_t^2 \lambda_{s+t} a_s \overline{a_{s+2t}}, \\ C_6 &= \sum_{s,t} a_s^2 a_t^2 \overline{a_{s+t}}^2. \end{aligned}$$

These are real: for instance, simultaneous negation of all frequency variables conjugates the summands. There is a real homogeneous form V_7 of degree seven for which

$$Q(\delta + f) = \delta^8 + 12\delta^4 S + 8\delta^3 B_5 + \delta^2(12A_6 + 12B_6 + 4C_6) + 8\delta V_7 + u^8. \quad (\text{F4})$$

Moreover,

$$|B_6|, |C_6| \leq A_6, \quad |V_7| \leq u^4 \sqrt{A_6}. \quad (\text{F5})$$

Here are details of the finite cube calculation in (F4). Terms with one, two, or three selected vertices vanish by centering. Of the 70 four-vertex supports, twelve are parallelograms and contribute S ; the other supports have independent forms. Of the 56 five-vertex supports, eight have no dependent four-subset. These are the cube-symmetric copies of $\{000, 100, 010, 001, 111\}$ and contribute B_5 ; every other five-vertex support has a Fourier constraint forcing a selected coefficient to be a_0 . All nonzero minors needed here are units in characteristic five. More explicitly, each of the twelve parallelograms has four extensions to a five-set. No five-set contains two different parallelograms: their affine planes would share three cube vertices, whereas two distinct planes meet in a line and a line contains at most two Boolean vertices. These 48 extensions have an independent fifth form, whose centered factor averages to zero. The other eight supports are the stated symmetric copies. A six-vertex support is classified by the Hamming distance of the two omitted vertices. Distances one, two, and three occur 12, 12, 4 times. For the representatives omitting respectively

$$\{000, 100\}, \quad \{000, 110\}, \quad \{000, 111\},$$

the Fourier constraints have solutions, in the lexicographic order of the remaining vertices,

$$(s+t, -s, -t, -s-t, s, t), \quad (t, s+t, -s-2t, -s-t, s, t), \quad (t, s, -s-t, -s-t, s, t).$$

Their products are exactly the summands defining A_6, B_6, C_6 . Cube symmetries preserve normalized averages by invertible changes of the cube parameters. The eight seven-vertex supports are also symmetric.

The bound for C_6 follows by taking absolute values. For B_6 , put $r = s + t$ and apply Cauchy–Schwarz:

$$\begin{aligned} |B_6| &\leq \sum_{r,t} \lambda_r \lambda_t |a_{r-t}| |a_{r+t}| \\ &\leq \left(\sum_{r,t} \lambda_r \lambda_t \lambda_{r-t} \right)^{1/2} \left(\sum_{r,t} \lambda_r \lambda_t \lambda_{r+t} \right)^{1/2} = A_6. \end{aligned}$$

For the last equality use $\lambda_{-t} = \lambda_t$.

For the septic bound, let $P(s, t)$ be the coefficient of δ when the Fourier coefficient $a_0 = 0$ in $T(s, t)$ is replaced by δ . It is the sum of four cubic expressions, one from each possible zero position in the four factors. Each expression has squared $\ell^2(s, t)$ norm A_6 ; for example, the first is $\overline{a_s} \overline{a_t} a_{s+t}$. The others are obtained by invertible changes of (s, t) and conjugation. Consequently $\|P\|_2 \leq 4\sqrt{A_6}$. The coefficient of δ in $\sum_{s,t} |T_{\delta+f}(s, t)|^2$ is $2 \operatorname{Re} \langle T, P \rangle = 8V_7$. Equation (F3) and Cauchy–Schwarz prove the septic bound.

III.58.2 A Fourier concentration penalty specific to order five

Choose a largest conjugate pair $\{\xi, -\xi\}$, and write

$$\rho = |a_\xi|, \quad b = |a_{2\xi}|, \quad D = \sum_{r \notin \{\xi, -\xi\}} |a_r|^4, \quad S = 2\rho^4 + D.$$

Since every nonzero character has order five, the doubled pair is distinct from the largest pair, and $2b^4 \leq D$. Define

$$E = \sum_{r+s=t+w} \lambda_r \lambda_s \lambda_t \lambda_w = \|\lambda * \lambda\|_2^2, \quad M = \sum_r |a_r|^8.$$

The slots $s = 0$ or $t = 0$ in (F3) contribute exactly $2E - S^2$. In the nonnegative energy sum retain the three families

$$(r = t, s = w), \quad (r = w, s = t), \quad (s = -r, w = -t).$$

Each has weight S^2 , each pairwise intersection has weight M , and their triple intersection has zero weight, since it forces $2r = 0$ and $a_0 = 0$. Thus $E \geq 3S^2 - 3M$. Conjugate pairing gives $M \leq S^2/2$, so

$$u^8 \geq 5S^2 - 6M \geq 2S^2. \quad (\text{F6})$$

Maximality of ρ also gives $M \leq \rho^4 S$ and hence

$$u^8 \geq (2 + 3D/S)S^2. \quad (\text{F7})$$

Order five supplies eight additional energy quadruples. Start with $(\xi, \xi, -\xi, -2\xi)$ and apply interchange of the two sides, interchange of the unequal entries on a side, and simultaneous negation.

These are eight distinct solutions of $r + s = t + w$, since $5\xi = 0$. They belong to none of the three preceding families, and each has weight $\rho^6 b^2$. The tail eighth-power mass is at most $D^2/2$, so retaining both the new quadruples and the concentration defect proves

$$\boxed{u^8 \geq 2S^2 + 12\rho^4 D + 16\rho^6 b^2.} \quad (\text{F8})$$

Indeed, the axes are at least $5S^2 - 6(2\rho^8 + D^2/2) + 16\rho^6 b^2 = 2S^2 + 12\rho^4 D + 16\rho^6 b^2$.

There is also a useful dimension-free estimate for the sextic quantity:

$$A_6 \leq 3\sqrt{E}\sqrt{D}, \quad E \leq \frac{u^8 + S^2}{2} \leq \frac{3}{4}u^8. \quad (\text{F9})$$

To prove the first inequality, set $q_r = \lambda_r$ outside $\{\xi, -\xi\}$ and zero on that pair. No three members of $\{\xi, -\xi\}$ sum to zero, because ξ has order five. Every nonzero term of A_6 therefore uses at least one tail index. The union bound, evenness of λ , and Cauchy–Schwarz give

$$A_6 \leq 3\langle q, \lambda * \lambda \rangle \leq 3\|q\|_2 \|\lambda * \lambda\|_2 = 3\sqrt{D}\sqrt{E}.$$

The second inequality follows from the axes contribution and (F6). This step is crucial: small fourth-power Fourier tail mass need not imply small L^2 mass, but it does control these particular incomplete-cube terms.

III.58.3 Uniform optimization and proof of the upper bound

The zero function is immediate. Otherwise write $g = f/u$, $t = u/\delta$, and use all Fourier letters in this subsection for g , so $\|g\|_{U^3} = 1$. Put

$$S_0 = 1/\sqrt{2}, \quad \eta = S_0 - S \geq 0, \quad F = Q(1 + tg) - 1 - 6\sqrt{2}t^4.$$

We may assume $0 < t \leq 1$ and $F > 0$; otherwise the assertion is immediate. The ordinary cube Cauchy–Schwarz bounds on the 28 sextic and eight septic terms give

$$F = -12\eta t^4 + 8t^5 B_5 + R, \quad |R| \leq 37t^6. \quad (\text{F10})$$

Since $|B_5| \leq \rho S < 1$, first $\eta = O(t)$. For small absolute t , $S \geq S_0/2$, and (F7) gives $D \leq 8\eta/3$. Consequently $\rho^4 = (S_0 - \eta - D)/2$ is bounded away from zero. Equation (F8) gives $b^2 = O(\eta)$.

If $a_\xi = \rho e^{i\theta}$, the two selected terms of B_5 are

$$2\rho^4 \operatorname{Re}(a_{2\xi} e^{-2i\theta}).$$

The remaining terms have total absolute value at most ρD . Thus $|B_5| = O(\sqrt{\eta} + \eta)$. Inserting this in (F10), absorbing the $O(t\eta)$ term, and solving the quadratic inequality in $\sqrt{\eta}$ yields

$$\eta = O(t^2), \quad D = O(t^2), \quad b = O(t). \quad (\text{F11})$$

All constants here are absolute.

Let $A_0 = 2^{-3/2}$. Since $\rho^4 = A_0 - (\eta + D)/2$, expanding (F8) using (F11) gives

$$\eta \geq \frac{3}{2}D + 2^{1/4}b^2 - O(t^4). \quad (\text{F12})$$

For clarity, $12\rho^4 D = 3\sqrt{2}D + O(t^4)$, $16\rho^6 b^2 = 2^{7/4}b^2 + O(t^4)$, and $1 - 2S^2 = 2\sqrt{2}\eta - 2\eta^2$; division by $2\sqrt{2}$ gives (F12).

Also $B_5 \leq S_0 b + D$. By (F9), $A_6 \leq 3\sqrt{D}$. Using the exact expansion (F4) and the sharper incomplete-cube bounds (F5) now gives

$$F \leq -18Dt^4 - 122^{1/4}b^2t^4 + 4\sqrt{2}bt^5 + 8Dt^5 \\ + 84t^6\sqrt{D} + 8\sqrt{3}t^7D^{1/4} + O(t^8).$$

For $t \leq 1$, the two terms linear in D are at most $-10Dt^4$. Completing the square in b yields

$$-122^{1/4}b^2t^4 + 4\sqrt{2}bt^5 \leq \frac{2^{3/4}}{3}t^6, \quad b = \kappa t \text{ at equality}, \quad \kappa = \frac{2^{1/4}}{6} = \frac{2^{5/4}}{12}.$$

Finally, on putting $y = D^{1/4}/t$, the remaining terms involving D are

$$t^8(-10y^4 + 84y^2 + 8\sqrt{3}y) = O(t^8),$$

because this polynomial is bounded above on $y \geq 0$. Thus $F \leq (2^{3/4}/3)t^6 + O(t^8)$ with absolute constants. Multiplication by δ^8 proves (F1).

III.58.4 An exact five-point model

On $\mathbb{Z}/5\mathbb{Z}$, set

$$a_0 = 0, \quad a_1 = \rho e^{i\pi/10}, \quad a_2 = b e^{i\pi/5}, \quad a_{-j} = \overline{a_j}.$$

The resulting function f is real and centered. Direct use of (F3) gives the exact identities

$$\|f\|_{U^3}^8 = 8\rho^8 + 16\rho^6 b^2 + 48\rho^4 b^4 + 16\rho^2 b^6 + 8b^8, \\ Q(\delta + f) = \delta^8 + 24\delta^4(\rho^4 + b^4) + 16\delta^3\rho^4 b + 96\delta^2\rho^4 b^2 + 48\delta^2\rho^2 b^4 + 32\delta\rho^4 b^3 + \|f\|_{U^3}^8. \quad (\text{F13})$$

A compact certificate is the following table for $T_{\delta+f}(s, t)$. Sign changes in either coordinate and interchange of the coordinates give the stated multiplicities.

(s, t)	multiplicity	$T_{\delta+f}(s, t)$
$(0, 0)$	1	$\delta^4 + 2\rho^4 + 2b^4$
$(0, 1)$	4	$2\delta^2\rho^2 + 2\rho^2 b^2 + b^4$
$(0, 2)$	4	$2\delta^2 b^2 + 2\rho^2 b^2 + \rho^4$
$(1, 1)$	4	$\delta^2\rho^2 + 2\delta\rho^2 b$
$(2, 2)$	4	$\delta^2 b^2$
$(1, 2)$	8	$\rho^2 b^2 + 2\delta\rho^2 b$

The multiplicities sum to 25. Each entry is a sum over five bases r ; the potentially additional terms cancel since $\cos(5\pi/10) = 0$. Squaring and summing proves (F13).

Choose $b = \kappa u^2/\delta$ and the unique positive ρ making the first expression in (F13) equal to u^8 . Such a ρ exists for sufficiently small u/δ , since the expression is strictly increasing in ρ and $8b^8 < u^8$. With $t = u/\delta$,

$$\rho^4 = u^4 \left(2^{-3/2} - 2^{-3/4}\kappa^2 t^2 + O(t^4) \right).$$

Substitution gives

$$Q(\delta + f) = \delta^8 + 6\sqrt{2}\delta^4 u^4 + \frac{2^{3/4}}{3}\delta^2 u^6 + O(u^8).$$

The weights lie in $[0, 1]$ when u is sufficiently small depending on δ , because $|f| \leq 2\rho + 2b \rightarrow 0$. Every nontrivial exponent-five group has a surjective character quotient onto $\mathbb{Z}/5\mathbb{Z}$; pullback along that quotient preserves normalized cube averages, means, and uniformity norms. This proves the lower bound on each fixed group. For a norm budget, the upper bound applies at the actual norm, and the positive terms in (F1) are increasing in that norm. This completes the proof of Theorem III.58.1.

III.58.5 Phase selection for families attaining the secondary term

Theorem III.58.3 (Necessary Fourier geometry). *Fix $\delta > 0$. Suppose f_j are real centered functions on finite groups of exponent five, $u_j = \|f_j\|_{U^3} \downarrow 0$, and*

$$Q(\delta + f_j) = \delta^8 + 6\sqrt{2}\delta^4 u_j^4 + \frac{2^{3/4}}{3} \delta^2 u_j^6 + o_\delta(u_j^6).$$

Choose a largest pair and write $a_{j,\xi_j} = \rho_j e^{i\theta_j}$ and $a_{j,2\xi_j} = b_j e^{i\psi_j}$; the latter is nonzero for large j . Then

$$\begin{aligned} b_j &\sim \kappa \delta^{-1} u_j^2, \\ e^{i(\psi_j - 2\theta_j)} &\longrightarrow 1, \quad \cos(5\theta_j) \longrightarrow 0, \\ \sum_{r \notin \{\xi_j, -\xi_j\}} |a_{j,r}|^4 &= o_\delta(u_j^6), \\ \rho_j^4 &= 2^{-3/2} u_j^4 - 2^{-3/4} \kappa^2 \delta^{-2} u_j^6 + o_\delta(u_j^6). \end{aligned}$$

These are necessary conditions, not a claimed sufficient characterization of all extremizing families. No L^2 approximation is asserted.

Proof. Work in the normalized variables of the upper proof and set $\phi = \psi - 2\theta$. Its estimates, retaining the actual phase, give

$$\frac{F}{t^6} \leq -18 \frac{D}{t^2} - 12 \cdot 2^{1/4} \left(\frac{b}{t}\right)^2 + 4\sqrt{2} \frac{b}{t} \cos \phi + o(1).$$

Here $D = O(t^2)$, $b = O(t)$, the sextic error divided by t^6 is $O(t)$, and the septic error is $o(1)$. The displayed expression has unique maximum $2^{3/4}/3$ in the variables $b/t \geq 0$, $D/t^2 \geq 0$ and $\cos \phi \leq 1$, attained when $b/t = \kappa$, $D/t^2 = 0$ and $\cos \phi = 1$. It follows that these three limits hold. The exact centered expansion then gives $\eta/t^2 \rightarrow 2^{1/4} \kappa^2$ and the claimed expansion of ρ^4 .

To detect the phase of the largest coefficient, also retain the four off-axis slots $(s, t) = (\pm 2\xi, \pm 2\xi)$. They have the same modulus. The contribution to $T(2\xi, 2\xi)$ from $r \in \langle \xi \rangle$ is exactly $2\rho^3 b \cos(3\theta + \psi)$. All remaining terms stay outside $\langle \xi \rangle$ in all four positions, so their sum has absolute value at most $D_{\text{out}} = \sum_{r \notin \langle \xi \rangle} |a_r|^4 \leq D$ by Hölder. Since $D = o(t^2)$ and $b \sim \kappa t$, the four slots supply

$$16\rho^6 b^2 \cos^2(3\theta + \psi) + o(t^2)$$

in addition to the axes lower bound. Applying the already established limits in (F8) forces $\cos(3\theta + \psi) \rightarrow 0$. Since $e^{i(\psi - 2\theta)} \rightarrow 1$, this is equivalent to $\cos(5\theta) \rightarrow 0$. Rescaling proves every stated conclusion. $\square$

III.58.6 Sharpness for subsets in growing dimension

Proposition III.58.4 (Indicator realization). *For every $0 < \delta < 1$, there are dimensions $n_j \rightarrow \infty$ and subsets $A_j \subset \mathbb{F}_5^{n_j}$, with densities $\delta_j \rightarrow \delta$ and $v_j = \|1_{A_j} - \delta_j\|_{U^3} \downarrow 0$, such that*

$$\frac{Q(1_{A_j}) - \delta_j^8 - 6\sqrt{2}\delta_j^4 v_j^4}{\delta_j^2 v_j^6} \rightarrow \frac{2^{3/4}}{3}.$$

Proof. Take the five-point model at parameters u_j decreasing geometrically and pull it back to $\mathbb{F}_5^n$ using the first coordinate. Include each point independently with its prescribed weight. With $N = 5^n$, at most $28/N$ of normalized cube parameter tuples have two coincident vertices: each equality of two distinct cube forms holds with probability $1/N$. On all other tuples, independence makes the expected indicator product equal the weighted product. The same holds for products centered at the fixed mean δ . Both expectation errors are therefore $O(1/N)$.

Changing one inclusion decision changes either cube polynomial by at most $8/N$, because each of the eight vertex forms has uniform fibers. Their variances are $O(1/N)$ by the bounded-difference variance estimate; the density variance is at most $1/(4N)$. Choosing n sufficiently large gives a realization in which the density, the indicator cube count, and the cube count centered at δ all approximate their weighted values to any prescribed error. Recentering at the realized density costs at most eight times the density error, by telescoping the bounded products.

Choose all these errors $o(u_j^{10})$. Then $v_j^8 = u_j^8 + o(u_j^{10})$, hence $v_j^4 = u_j^4 + o(u_j^6)$ and $v_j/u_j \rightarrow 1$. The model expansion transfers to the displayed ratio. The geometric choice of u_j makes the actual v_j strictly decrease after finitely many terms are discarded. $\square$

III.59 A closer expansion of the five-point example

Source 23, delivered Appendix A.

This appendix records a lower-model computation useful for the second research question. It does not identify the global tertiary optimum. Use the model of (F13), with $\kappa = 2^{1/4}/6$, $t = u/\delta$, and $b = \kappa u^2/\delta$. Let $A = 2^{-3/2}$. Solving the norm equation to two further orders gives

$$\frac{\rho^4}{u^4} = A - A^{1/2}\kappa^2 t^2 - 2\kappa^4 t^4 + O(t^6).$$

To verify this, substitute $A + at^2 + ct^4$ into the first equation in (F13), divide by u^8 , and compare coefficients. The constant coefficient is $8A^2 = 1$; the next two equations are

$$16Aa + 16A^{3/2}\kappa^2 = 0,$$

$$8a^2 + 16Ac + 24A^{1/2}a\kappa^2 + 48A\kappa^4 = 0.$$

They give $a = -A^{1/2}\kappa^2$ and $c = -2\kappa^4$. The positive root is analytic in t^2 near zero, since the derivative of the constant norm equation with respect to ρ^4/u^4 is $16A \neq 0$, so the expansion has the indicated remainder.

Substitution into the weighted-count formula gives the coefficient of u^8 as

$$-24\kappa^4 - 16A^{1/2}\kappa^3 + 96A\kappa^2 + 1 = -\frac{1}{27} - \frac{2}{27} + \frac{4}{3} + 1 = \frac{20}{9}.$$

Consequently this particular family satisfies

$$Q(\delta + f) = \delta^8 + 6\sqrt{2}\delta^4 u^4 + \frac{2^{3/4}}{3}\delta^2 u^6 + \frac{20}{9}u^8 + O(u^{10}/\delta^2). \quad (\text{III.59.1})$$

It follows that a global tertiary coefficient, if it exists in this normalization, must be at least $20/9$. The uniform upper argument in Section III.58 bounds that coefficient but does not equate it with the model value.

III.60 Source 22’s chapter: degree filtrations of polynomial colourings [write]

Added in the fifth write, 6 October 2026.

The next five sections are source 22’s Sections 6–10, on colourings $Q : \mathbb{F}_q^n \rightarrow \mathbb{F}_q^r$ by polynomials of bounded degree and the directions h for which a prescribed pattern of equal colours occurs from a base point. *Logically the chapter continues source 10’s Section III.30* (the dimension threshold for symmetric progressions) *and source 14’s Section III.36* (exact supersaturation in low degree), which source 22 cites as prior project work and whose ingredients it reproves. It is printed here, at the end of Part III, because no number printed earlier may move. Source 22’s introduction is Section F.27; its Sections 2–5 are in Part II and Sections 11–12 in Part V.

What is proved, in one paragraph. For every prescribed equality pattern and degree caps, the minimum number of successful directions at a prescribed centre is $q^{\max(n-B,0)}$, with B the maximum-weight column-basis sum of an explicit evaluation matrix (Theorem III.61.5); for symmetric progressions this removes the degree restriction of source 14’s Theorem III.36.1 (Corollary III.62.2), and a monochromatic analogue holds in every degree. The scalar quintic and sextic pointwise budget for symmetric four-term progressions is exactly eight. Further: all scalar quadratic minimizers (Theorem III.63.1), no independent additive extension of the cubic norm core (Theorem III.64.2) with an exact defect count (Theorem III.64.3), zero-defect rigidity (Corollary III.64.4), and a quantitative excess for norm layers of degree divisible by the characteristic (Theorem III.65.3). Warning’s second theorem, the Delsarte–Goethals–MacWilliams rigidity and the norm constructions of Shi and Dong are credited as prior work. The global quintic avoidance range $4 \leq d \leq 8$ stays open (item III.19); the value 7,810,000 is a lower bound, not a count.

Relation to the formal statements. The colouring problem has no Lean statement; nothing in this chapter is formalized.

III.61 Polynomial zero sets and prescribed equality patterns

Source 22, delivered Section 6.

From this section onward, q denotes a prime power and $\mathbb{F}_q$ its field; p is its characteristic, and n is the dimension of a coloring domain. This avoids confusing field size with the $2d$ cubes of the preceding sections.

III.61.1 The classical zero-counting input

We use a familiar polynomial support bound. The proof is included because the distinction between a lower bound and its equality case is useful here.

Lemma III.61.1 (Minimum support of a polynomial function). *Let F be a nonzero polynomial function on $\mathbb{F}_q^n$, represented by a polynomial of total degree at most $R < n(q-1)$. Write $R = a(q-1) + b$, $0 \leq b < q-1$. Then*

$$|\text{supp } F| \geq (q-b)q^{n-a-1}.$$

In particular, degree at most $B(q-1)$ and nonempty support imply $|\text{supp } F| \geq q^{n-B}$ for $0 \leq B \leq n$.

Proof. Reduce each individual exponent using $X_i^q = X_i$. Reduction does not increase total degree, and produces the unique polynomial of individual degrees at most $q-1$ representing the function. Uniqueness follows by induction from the univariate root bound.

Successively take a leading coefficient in one variable. If the degree in that variable is e , then above each input where its leading coefficient is nonzero, at least $q-e$ points are in the support. Induction therefore produces exponents $0 \leq e_i \leq q-1$ of an actual monomial, with $\sum e_i \leq R$, such that $|\text{supp } F| \geq \prod_i (q-e_i)$. For a fixed sum, this product is minimized by making the exponents as unequal as possible. Indeed, for $0 < e_i \leq e_j < q-1$, replacing the pair by e_i-1, e_j+1 decreases its product. The minimum has a exponents $q-1$, one exponent b , and the others zero, giving the stated bound. The case $B = n$ merely says that the support is nonempty. $\square$

Corollary III.61.2 (Warning's second bound). *Suppose $f_1, \dots, f_s \in \mathbb{F}_q[X_1, \dots, X_n]$ have sum of degrees at most B and have a common zero. Their common zero set has at least $q^{\max(n-B, 0)}$ elements.*

Proof. Its indicator is represented by $\prod_i (1 - f_i^{q-1})$, a polynomial of degree at most $B(q-1)$. Apply Lemma III.61.1 when $B \leq n$, and otherwise use the given zero. $\square$

These are classical results, with Warning's original theorem and an elementary modern proof in [40, 56]. The following equality case is the relevant part of the minimum-weight generalized Reed–Muller theorem [96, 57].

Lemma III.61.3 (Rigidity at minimum support). *If a nonzero polynomial function of degree at most $B(q-1)$ has support of size exactly q^{n-B} , where $0 \leq B \leq n$, its support is an affine subspace of dimension $n-B$.*

Proof. The cases $B = 0, n$ are immediate. Put $e = n - B \geq 1$ and write S for the support. Consider a family of parallel affine hyperplanes and, after an affine change of variables, write it as $X_1 = c$. Let t members miss S . If $1 \leq t \leq q-2$, the reduced polynomial factors as

$$F(X) = \prod_{c: S \cap \{X_1=c\} = \emptyset} (X_1 - c)G(X), \quad \deg G \leq B(q-1) - t.$$

The factorization follows because a reduced restriction vanishing on an entire field grid is the zero polynomial. On every remaining hyperplane, Lemma III.61.1, in $n-1$ variables, gives at least $(t+1)q^{e-1}$ support points: the residual degree is $(B-1)(q-1) + (q-1-t)$. Therefore

$$|S| \geq (q-t)(t+1)q^{e-1} > q^e,$$

a contradiction. Thus every parallel family misses either zero or $q-1$ members. In a family missing none, every member has exactly q^{e-1} support points, since each has at least this many and the total is q^e .

Some family must miss a member. Otherwise every affine hyperplane contains exactly q^{e-1} points of S . Choose $x \notin S$ and count pairs consisting of a hyperplane through x and a point of S on it. Counting first hyperplanes and then points gives

$$q^{e-1}(q^n - 1) = q^e(q^{n-1} - 1),$$

which is impossible. Consequently S lies in one affine hyperplane. Move it to $X_1 = 0$. Uniqueness of reduced representatives gives

$$F(X_1, X') = (1 - X_1^{q-1})G(X').$$

Thus $\deg G \leq (B-1)(q-1)$, and its support still has size q^e . Induction on B proves that this support, and hence S , is an affine e -space. $\square$

III.61.2 Norm forms and their fibers

For each positive integer j , choose a basis of $\mathbb{F}_{q^j}$ over $\mathbb{F}_q$. Its field norm N_j is a homogeneous polynomial of degree j on a j -dimensional vector space. It vanishes exactly at zero. On nonzero elements,

$$N_j(x) = x^{(q^j-1)/(q-1)}.$$

The multiplicative group of a finite field is cyclic, so the norm is onto $\mathbb{F}_q^*$ and every nonzero fiber has exactly

$$\frac{q^j - 1}{q - 1} \tag{III.61.1}$$

elements; the zero fiber has one. One can also define the norm as the determinant of multiplication by x , making its homogeneous degree immediate. We use these standard finite-field facts as algebraic background, not as new constructions. Layered norms are used in the related progression-coloring work of Shi and Dong [33].

III.61.3 The evaluation matrix

Choose distinct nodes $\alpha_1, \dots, \alpha_\ell \in \mathbb{F}_q$ and a partition Π of their indices. The desired pattern says that the colors of $z + \alpha_i h$ agree whenever i, j belong to one block of Π . Choose a spanning tree on each block and write $\mathcal{E}$ for the resulting set of ordered edges. Equality on those edges is equivalent to equality on each block; the choice of trees will not affect the resulting row space.

For a scalar polynomial of degree at most D , define columns

$$v_j = (\alpha_a^j - \alpha_b^j)_{(a,b) \in \mathcal{E}}, \quad 1 \leq j \leq D. \tag{III.61.2}$$

Scan them from degree D down to degree one. Retain j if v_j does not belong to the span of previously retained columns. Let $J(\Pi, D)$ be the retained indices, and put

$$B(\Pi, D) = \sum_{j \in J(\Pi, D)} j. \tag{III.61.3}$$

For an empty edge set the span and budget are zero. A specified sparse set of allowed degrees is handled by scanning only that set.

Lemma III.61.4 (Optimal constant recombination by degree). *Give independent formal symbols C_j weight j . The vector relation $\sum_j v_j C_j = 0$ has an equivalent basis of scalar equations whose degrees are exactly the retained indices in $J(\Pi, D)$. No equivalent basis formed by constant linear combinations has a smaller sorted degree list or a smaller sum of degrees. Equivalently, $B(\Pi, D)$ is the maximum sum of degree indices of a column basis.*

Proof. For every retained b , take the dual coordinate functional λ_b on the span of the retained basis. A rejected column of degree j is a combination of retained columns of degrees larger than j . Therefore

$$\lambda_b \left(\sum_j v_j C_j \right) = C_b + \sum_{j < b} a_{b,j} C_j.$$

The equations have degrees b and are equivalent to the vector relation. For optimality, the functionals giving degree at most t form the annihilator of $\text{span}\{v_j : j > t\}$. Its dimension is $r - \text{rank}(v_{t+1}, \dots, v_D)$, where $r = \text{rank}(v_1, \dots, v_D)$. No equation basis can have more than that number of members of degree at most t . Greedy elimination attains these dimensions at every t , proving the degree comparison. The usual descending exchange argument, or the same ranks of suffixes, identifies its column basis as one of maximum total weight. $\square$

Theorem III.61.5 (Exact fixed-center pattern minimum). *Let $Q = (Q_1, \dots, Q_r) : \mathbb{F}_q^n \rightarrow \mathbb{F}_q^r$ have coordinate degree caps $D_1, \dots, D_r$, and let*

$$S_z(Q) = \{h : Q(z + \alpha_a h) = Q(z + \alpha_b h) \text{ for every } (a, b) \in \mathcal{E}\}.$$

Set $B = \sum_{a=1}^r B(\Pi, D_a)$. For every prescribed z ,

$$\boxed{\min_Q |S_z(Q)| = q^{\max(n-B, 0)}}. \quad (\text{III.61.4})$$

If $n \geq B$ and equality holds, $S_z(Q)$ is an $\mathbb{F}_q$ -linear subspace of dimension $n - B$. The theorem also holds with specified sparse degree supports at the prescribed center, using their greedy budgets.

Proof. Expand formally $Q_a(z + TH) = \sum_{j=0}^{D_a} C_{a,j,z}(H) T^j$. Every $C_{a,j,z}$ is homogeneous of degree j in H , or zero. The equality pattern is precisely $\sum_j v_j C_{a,j,z} = 0$ in each output coordinate. Lemma III.61.4 converts these to equations with total degree sum at most B . They vanish at zero. Corollary III.61.2 proves the lower bound. Their common-zero indicator has degree at most $B(q-1)$, so Lemma III.61.3 proves the equality assertion; the affine subspace contains zero and is therefore linear.

For sharpness in dimension B , allocate a disjoint block of dimension j to every retained degree j in every output, and make that output the sum of its block norms. At the origin its pattern vector is $\sum_{j \in J} v_j N_j(h^{(j)})$. Independence of the retained columns forces every norm, hence every block, to vanish. Thus $S_0 = \{0\}$. For $n > B$ add ignored coordinates; for $n < B$ restrict along a linear injection into the B -dimensional construction. Translation places the sharp example at any prescribed z . The case $B = 0$ is immediate. The same proof, with only the allowed columns, gives the sparse version. $\square$

Remark III.61.6 (What is optimized). Lemma III.61.4 concerns constant linear recombinations of the evaluation equations. Polynomial multipliers could change their degree presentation for a particular coloring. The stronger extremal conclusion (III.61.4) does not depend on limiting proof methods: the norm example attains the bound among all admissible colorings at the prescribed center.

III.62 Arithmetic-progression specializations

Source 22, delivered Section 7.

III.62.1 Monochromatic progressions in every degree

Take nodes $0, 1, \dots, \ell - 1$ and require one common color, with $p > \ell$. Comparing each node to zero gives $v_j = (1^j, 2^j, \dots, (\ell - 1)^j)^\top$. Any consecutive $\ell - 1$ columns are independent: factor nonzero powers out of the rows and use the Vandermonde determinant. If $D < \ell - 1$, the first D columns are independent as well. Consequently, for $r_D = \min(D, \ell - 1)$,

$$B_\ell^{\text{mono}}(D) = \frac{r_D(2D - r_D + 1)}{2} = \begin{cases} D(D + 1)/2, & D \leq \ell - 1, \\ (\ell - 1)(2D - \ell + 2)/2, & D \geq \ell - 1. \end{cases} \quad (\text{III.62.1})$$

This is an exact pointwise budget in the sense of Theorem III.61.5. Unlike symmetry alone, the required pattern is genuinely monochromatic and thus directly matches the type of configuration in Szemerédi's theorem. Its finite-field, bounded-polynomial-coloring scope remains essential.

Corollary III.62.1 (Exact global minimum in the interpolation range). *Suppose all degree caps satisfy $D_a \leq \ell - 1$, and put $B = \sum_a D_a(D_a + 1)/2$. The minimum number of ordered nontrivial monochromatic ℓ -term progressions over all such colorings is*

$$\min_Q A_\ell^{\text{mono}}(Q) = q^n \left(q^{\max(n-B, 0)} - 1 \right). \quad (\text{III.62.2})$$

One polynomial map attains the minimum at every base point.

Proof. Sum the pointwise lower bound minus its zero direction over all q^n base points. For the matching example, in output a use independent norm blocks of every degree $1, \dots, D_a$. If a direction is nonzero on the controlled coordinates, choose its highest nonzero block in some output. The polynomial obtained on that affine line has a nonzero leading coefficient equal to the corresponding norm. It has degree at most $D_a < \ell$, so it cannot have one value at all ℓ distinct nodes. Thus only ignored directions succeed at every base point. Restriction to a linear subspace gives the case $n < B$. $\square$

III.62.2 Symmetric progressions without a degree restriction

Let $\ell = 2m \geq 4$, $p > \ell$, and use the centered nodes $\pm\alpha_i$, where $\alpha_i = i - \frac{1}{2}$, $1 \leq i \leq m$. The symmetry condition is

$$Q(z + \alpha_i h) = Q(z - \alpha_i h) \quad (1 \leq i \leq m).$$

For four terms it permits a color sequence (a, b, b, a) with $a \neq b$; it must not be identified with the monochromatic condition.

Only odd coefficients survive reflection. Dividing the i th equation by $2\alpha_i$, its odd-degree columns are

$$((\alpha_i^2)^{j-1})_{i=1}^m, \quad j = 1, \dots, s, \quad s = \lceil D/2 \rceil.$$

The squares α_i^2 are distinct and nonzero, since their differences are $(i - j)(i + j - 1)$ and $p > \ell$. The last $\min(m, s)$ columns are therefore independent. We obtain the following formula.

Corollary III.62.2 (Exact arbitrary-degree symmetric budget). *For output degree caps D_a , set $s_a = \lceil D_a/2 \rceil$, $t_a = \min(m, s_a)$ and*

$$B_{\text{sym}} = \sum_a t_a(2s_a - t_a) = \sum_a (s_a^2 - (s_a - m)_+^2). \quad (\text{III.62.3})$$

At every prescribed center,

$$\min_Q |S_z(Q)| = q^{\max(n - B_{\text{sym}}, 0)}.$$

If $n \geq B_{\text{sym}}$, equality forces an $\mathbb{F}_q$ -linear space of directions of dimension $n - B_{\text{sym}}$. In all degrees,

$$A_\ell^{\text{sym}}(Q) \geq q^n (q^{\max(n - B_{\text{sym}}, 0)} - 1). \quad (\text{III.62.4})$$

When every $D_a \leq \ell$, simultaneous equality is attainable at all centers, so (III.62.4) is then the exact global minimum.

Proof. The pointwise statements follow from the evaluation columns and Theorem III.61.5. The retained odd degrees in an output are $2s_a - 2t_a + 1, \dots, 2s_a - 1$, whose sum is (III.62.3). The correspondence $z = a + (\ell - 1)h/2$ between starting-point and centered parametrizations is bijective, giving the global lower bound.

If $D_a \leq \ell$, use disjoint norm blocks of odd degrees $1, 3, \dots, 2s_a - 1$. Reflection at the m pairs forces every odd coefficient to vanish. The highest nonzero direction block would give a nonzero highest odd coefficient, a contradiction. Only ignored directions remain at every center, proving equality. This recovers the low-degree result of [106]; the arbitrary-degree pointwise statement is the extension proved here. $\square$

Degree cap D	1	2	3	4	5	6	7	8	9	10
$B_4^{\text{sym}}(D)$	1	1	4	4	8	8	12	12	16	16
$B_4^{\text{mono}}(D)$	1	3	6	9	12	15	18	21	24	27

III.62.3 The quintic obstruction is exactly pointwise

For $\ell = 4$ and $D = 5$ or 6 , the sharp budget is eight. In eight variables the polynomial

$$Q(u, v) = N_3(u) + N_5(v), \quad (u, v) \in \mathbb{F}_q^3 \oplus \mathbb{F}_q^5,$$

has $S_0(Q) = \{0\}$: its two odd evaluation columns of degrees three and five are independent. This is valid even in characteristic five; the relevant squared-node Vandermonde determinant remains nonzero.

Therefore the universal pointwise statement “every scalar quintic in eight variables has a nonzero symmetric direction at every prescribed center” is false. A proof that lowers the global avoidance bound from eight must use compatibility among different centers or another global feature. The previous cubic norm construction gives a globally avoiding example in dimension four. The interval between four and eight for the global quintic/sextic question is not closed by the present theorem.

For a sparse example, over $\mathbb{F}_5$ take allowed odd degrees $\{1, 5\}$. Fermat’s identity gives identical evaluation columns, so the greedy budget is five rather than six. For full quintic support $\{1, 3, 5\}$, the retained degrees are three and five, giving eight. This illustrates why the actual evaluation matrix, including characteristic-dependent relations, is more informative than a formal sum of all odd degrees.

III.63 Complete quadratic classification and counts

Source 22, delivered Section 8.

Let q be odd and let $P : V = \mathbb{F}_q^n \rightarrow \mathbb{F}_q$ have degree at most two. The structural statements below only require odd characteristic; for the count of distinct arithmetic progressions retain $p > \ell$ as above. Write uniquely

$$P(x) = B(x, x) + L(x) + c$$

with B a symmetric bilinear form and L a linear functional. Put $R = \text{rad}(B)$ and $r = \text{rank}(B)$. At every center z ,

$$P(z + \alpha h) - P(z - \alpha h) = 2\alpha(2B(z, h) + L(h)).$$

Thus any nonempty reflection configuration with nonzero nodes has

$$S_z(P) = \ker(2B(z, \cdot) + L).$$

Theorem III.63.1 (All scalar quadratic minimizers). *Assume $n \geq 1$. The following are equivalent:*

1. $|S_z(P)| = q^{n-1}$ for every z .
2. $L|_R \neq 0$.
3. Some $u \in V$ satisfies $P(x + su) = P(x) + s$ for all $x \in V$ and $s \in \mathbb{F}_q$.
4. An invertible affine change of source coordinates puts P in the form $t + B_0(w, w) + c'$, where B_0 is nondegenerate on an r -dimensional space, and the other $n - r - 1$ variables are unused.

If $L|_R = 0$, there are exactly q^{n-r} exceptional centers, at which $S_z = V$. At every other center, $|S_z| = q^{n-1}$. Hence in this case

$$A_\ell^{\text{sym}}(P) = q^n(q^{n-1} - 1) + q^{n-r}(q^n - q^{n-1}).$$

Among non-minimizing scalar quadratic maps, the least possible excess above the global minimum is $q^n - q^{n-1}$, attained precisely when B is nondegenerate.

Proof. The linear map $z \mapsto B(z, \cdot)$ has kernel R and image exactly $R^\perp \subseteq V^*$. Therefore $2B(z, \cdot) + L$ is never zero if and only if $L \notin R^\perp$, namely $L|_R \neq 0$. In that case choose $u \in R$ with $L(u) = 1$; the translation identity follows by expanding $P(x + su)$. Conversely the translation identity, first at $s = 1$ and varying x , forces $B(x, u) = 0$ for all x and then $L(u) = 1$.

Split $V = W \oplus R$, with $B|_W$ nondegenerate, and split $R = \mathbb{F}_q u \oplus (R \cap \ker L)$. Completing the square in W removes the linear term on W , giving the asserted normal form. The normal form immediately gives the minimal counts. If $L \in R^\perp$, the affine equation $2B(z, \cdot) = -L$ has exactly q^{n-r} solutions. A nonzero linear functional has kernel size q^{n-1} ; the zero functional has kernel size q^n . Summing the excess at the exceptional centers gives the count. Since $r \leq n$, its least positive value occurs exactly at $r = n$. $\square$

III.64 Why independent additive blocks cannot extend the norm coloring

Source 22, delivered Section 9.

In this section let $\ell = 2m$, $p > \ell$, and $\alpha_i = i - 1/2$ for $1 \leq i \leq m$.

The obstruction below is global and permits an arbitrary added function, of any polynomial degree or with no polynomial presentation imposed. It excludes a large and natural class of attempts to improve the m^2 -dimensional norm-layer coloring by appending new input coordinates.

Let

$$U = \bigoplus_{j=1}^m \mathbb{F}_q^{2j-1}, \quad C(u) = \sum_{j=1}^m N_{2j-1}(u_j), \quad \dim U = m^2.$$

Here each N_t is a degree- t field norm and N_1 is the identity. For a core center $a \in U$, define its reflection map

$$\Phi_a(h) = (C(a + \alpha_i h) - C(a - \alpha_i h))_{i=1}^m \in \mathbb{F}_q^m.$$

Lemma III.64.1 (Surjectivity of every core reflection map). *For every $a \in U$, the map $\Phi_a : U \rightarrow \mathbb{F}_q^m$ is onto.*

Proof. The odd part of $C(a + Th)$ has precisely the possible degrees $1, 3, \dots, 2m - 1$. The reflection evaluation matrix for those degrees is invertible, so any desired reflected-value vector is equivalent to prescribing these m odd coefficients. Choose directions block by block, starting with the largest block. The coefficient of T^{2m-1} is exactly $N_{2m-1}(h_m)$, which can be set arbitrarily because the norm map is onto $\mathbb{F}_q$. Once h_m is fixed, the coefficient of T^{2m-3} is $N_{2m-3}(h_{m-1})$ plus an already determined contribution from the larger block. Choose h_{m-1} to make this coefficient correct. Continue downward; at each stage smaller blocks cannot affect the current coefficient, and the norm of the current block can attain every desired value. The final linear block sets the coefficient of T . This realizes the prescribed reflection vector. $\square$

Theorem III.64.2 (No independent additive extension). *Let $W = \mathbb{F}_q^e$, and let $R : W \rightarrow \mathbb{F}_q$ be any function. Set*

$$P(u, w) = C(u) + R(w).$$

For every center (a, b) , the projection

$$S_{(a,b)}(P) \longrightarrow W, \quad (h, g) \longmapsto g,$$

is surjective. Consequently

$$|S_{(a,b)}(P)| \geq q^e, \quad A_\ell^{\text{sym}}(P) \geq q^{m^2+e}(q^e - 1).$$

If $e > 0$, the coloring cannot avoid nontrivial symmetrically colored ℓ -term progressions. This holds irrespective of the degree of R .

Proof. Fix an extra direction $g \in W$. The m reflected differences of $R(b + Tg)$ form a vector in $\mathbb{F}_q^m$. Surjectivity of Φ_a supplies a core direction h giving its negative. Thus (h, g) is a symmetric direction. For $g \neq 0$ it is nonzero. The projection conclusion gives the pointwise bound, and summation over the q^{m^2+e} centers gives the global count. $\square$

The same proof applies to several output coordinates, each with its own complete m -layer norm core, and an arbitrary added vector-valued function. For the scalar four-term problem, it rules out every construction of the form $x_0 + N_3(u) + R(w)$ with at least one added variable. In particular, the failure of adding independent higher norm blocks is not confined to the Artin–Schreier example in characteristic five.

III.64.1 Exact counts for the cubic core

For $\ell = 4$, the preceding lower bound has an exact defect formula. Let

$$P(a, u, w) = a + N_3(u) + R(w), \quad (a, u, w) \in \mathbb{F}_q \oplus \mathbb{F}_q^3 \oplus W.$$

For $b, g \in W$, put

$$D_R(b, g) = R(b + 3g/2) - 3R(b + g/2) + 3R(b - g/2) - R(b - 3g/2),$$

and let $T_b(R) = \#\{g \in W : D_R(b, g) \neq 0\}$. $D_R(b, g)$ is the third forward difference $\Delta_g^3 R(b - 3g/2)$.

Theorem III.64.3 (Exact added-block defect). *For every center (a, u, b) ,*

$$|S_{(a, u, b)}(P)| = q^e + (q^2 + q)T_b(R).$$

Consequently

$$A_4^{\text{sym}}(P) = q^{4+e}(q^e - 1) + q^4(q^2 + q) \sum_{b \in W} T_b(R).$$

Thus the universal added-block lower bound is attained if and only if every pure third additive difference of R vanishes.

Proof. Fix g . Let E_1 and E_3 denote the reflected differences of the cubic core at the parameters $1/2$ and $3/2$. For the cubic core, write its odd coefficients as C_1 and C_3 . Then $E_1 = C_1 + C_3/4$ and $E_3 = 3C_1 + 27C_3/4$. Therefore the two cancellation equations force

$$N_3(h_u) = C_3 = -D_R(b, g)/6.$$

Each choice of h_u satisfying this equation determines a unique linear direction h_a from the first reflection equation. The cubic norm has one preimage of zero and $q^2 + q + 1$ preimages of each nonzero value. Thus the fiber over g has size one if $D_R(b, g) = 0$ and size $q^2 + q + 1$ otherwise. Summing over g proves the pointwise formula. It is independent of the core center (a, u) , of which there are q^4 ; summing over all centers gives the stated global formula. $\square$

III.64.2 The zero-defect functions over the prime field

Corollary III.64.4 (Zero-defect rigidity). *If W is regarded as a vector space over its prime field $\mathbb{F}_p$ (with $p > 3$), then all pure third differences vanish if and only if R is a polynomial of total degree at most two in $\mathbb{F}_p$ -linear coordinates, with coefficients in $\mathbb{F}_q$. Equivalently, $R(x) = c + L(x) + B(x, x)$ with L $\mathbb{F}_p$ -linear and B symmetric $\mathbb{F}_p$ -bilinear. For $q = p$ these are precisely the usual scalar quadratic polynomials. For proper extension fields this includes, for example, additive Frobenius maps, and should not be replaced by an $\mathbb{F}_q$ -quadratic claim.*

Proof. Choose an $\mathbb{F}_p$ -basis of W and use the unique reduced polynomial representation $R(X) = \sum_a c_a X^a$, in which every $a_i < p$. If all pure third differences vanish, the restriction $t \mapsto R(x + th)$ on each prime-field line has reduced degree at most two: otherwise applying the step-one third difference to a polynomial of degree $j \in \{3, \dots, p-1\}$ leaves a nonzero leading coefficient $j(j-1)(j-2)$. The coefficient of t^3 in the reduced polynomial for $R(X + tH)$ therefore vanishes as a polynomial function of (X, H) . Its expansion contains terms

$$c_a \binom{a}{e} X^{a-e} H^e \quad \text{for } e \leq a \text{ and } |e| \equiv 3 \pmod{p-1}.$$

Every displayed exponent in each X_i and H_i is below p , and different pairs (a, e) yield different monomials. Hence there can be no cancellation between them. If $|a| \geq 3$, choose $e \leq a$ with $|e| = 3$; all the binomial coefficients are nonzero in characteristic p , forcing $c_a = 0$. Thus R has total degree at most two. The converse follows immediately from taking third differences. Polarization in odd characteristic gives the equivalent bilinear description. $\square$

[write] **6 October 2026; sources 34–39.** Source 38's Theorem III.88.2 gives the exact six-term count for the linear–cubic–quintic core with an independent addition (Section III.87).

III.65 A quantitative global obstruction to the high-degree norm construction

Source 22, delivered Section 10.

Let q be a power of an odd prime p . Write $E_t = \mathbb{F}_{q^t}$, regarded as a t -dimensional vector space over $\mathbb{F}_q$, and write $N_t = N_{E_t/\mathbb{F}_q}$. The function N_t is represented by a homogeneous polynomial of degree t in any linear coordinate basis. All progressions below are ordered start–direction pairs; directions are not identified up to scaling or orientation.

Lemma III.65.1 (Exact constant-line ratios for the degree- p norm). *Define*

$$\mathcal{T} = \{\theta \in E_p : \theta^q - \theta \in \mathbb{F}_p^*\}.$$

Then $|\mathcal{T}| = q(p-1)$, every element of $\mathcal{T}$ is nonzero, and for $\theta \in E_p$ the following are equivalent:

1. $N_p(\theta + j)$ is independent of $j \in \mathbb{F}_p$;
2. $\theta \in \mathcal{T}$;
3. as a polynomial in T ,

$$N_p(\theta + T) = T^p - T + \theta^p - \theta.$$

Consequently, the number of pairs $(a, h) \in E_p \times E_p^$ for which $N_p(a + jh)$ is constant for every $j \in \mathbb{F}_p$ is exactly*

$$(q^p - 1)q(p-1).$$

For each fixed $a \neq 0$, there are exactly $q(p-1)$ such directions; for $a = 0$, there are none.

Proof. For $c \in \mathbb{F}_p^*$ the polynomial $X^q - X - c$ has q distinct roots in an algebraic closure, since its derivative is -1 . If θ is a root, iteration gives $\theta^{q^j} = \theta + jc$, and hence $\theta^{q^p} = \theta$. Thus all these roots lie in E_p . The root sets for different c are disjoint, proving $|\mathcal{T}| = q(p-1)$. They avoid zero. For such a θ , its p Frobenius conjugates are exactly $\theta + \mathbb{F}_p$, because $c \neq 0$. Therefore

$$N_p(\theta + T) = \prod_{j \in \mathbb{F}_p} (\theta + T + j) = (\theta + T)^p - (\theta + T) = T^p - T + \theta^p - \theta.$$

This proves (ii) $\Rightarrow$ (iii) $\Rightarrow$ (i).

Conversely, put $P(T) = N_p(\theta + T)$, a monic polynomial of degree p over $\mathbb{F}_q$. If (i) holds, $P(T) - P(0)$ vanishes at all p elements of $\mathbb{F}_p$; it is monic of degree p , so $P(T) - P(0) = T^p - T$. In particular P is separable. If $\theta \in \mathbb{F}_q$ then $P(T) = (T + \theta)^p$, a contradiction. Since $[E_p : \mathbb{F}_q] = p$, the element θ therefore has degree p over $\mathbb{F}_q$. The roots of P are both its distinct Frobenius conjugates $-\theta^{q^j}$ and the translates $-\theta + \mathbb{F}_p$. Hence $\theta^q - \theta \in \mathbb{F}_p$; this difference is nonzero because $\theta \notin \mathbb{F}_q$. Thus (i) $\Rightarrow$ (ii).

For $h \neq 0$, multiplicativity of the norm shows that the progression condition is equivalent to $a/h \in \mathcal{T}$. There are $q^p - 1$ choices of h and $q(p-1)$ choices of this ratio. For fixed $a \neq 0$, the directions are $h = a/\theta$, $\theta \in \mathcal{T}$, and are distinct; $a = 0$ gives ratio zero and therefore no direction. $\square$

The word “exact” in this lemma concerns constancy on the full p -point prime-field progression. It does not assert an exact count for the weaker reflection equalities on $\ell < p$ points.

Proposition III.65.2 (Every norm block of degree divisible by the characteristic). *Suppose $p \mid t$. At every $a \in E_t^*$ there are at least $q(p-1)$ distinct $h \in E_t^*$ such that $N_t(a + jh)$ is constant for all $j \in \mathbb{F}_p$. Hence this norm has at least*

$$(q^t - 1)q(p-1)$$

ordered nonconstant-direction progressions on which its value is constant.

Proof. The subfield $E_p \subset E_t$ exists because $p \mid t$. Use $\mathcal{T} \subset E_p^*$ from the lemma. For $x \in E_p$, transitivity of the norm gives

$$N_t(x) = N_p(x)^{t/p}.$$

For each $a \neq 0$ and $\theta \in \mathcal{T}$, put $h = a/\theta$. Then

$$N_t(a + jh) = N_t(h)N_p(\theta + j)^{t/p},$$

which is independent of j . Distinct θ give distinct directions. Counting over the nonzero starts proves the assertion. $\square$

Theorem III.65.3 (Quantitative excess for a layered norm map). *Consider a direct-sum norm construction on B controlled coordinates, with any number $f \geq 0$ of additional ignored coordinates, so the domain has dimension $n = B + f$. Each controlled block has dimension t and contributes its norm N_t to one specified coordinate of the coloring. Different blocks use disjoint coordinates. Let $\mathcal{J}$ be any collection of its blocks whose degrees t_j are divisible by p . For any even $\ell < p$, let S_z denote the directions producing reflection symmetry about z on the standard centered ℓ -term progression, including zero. Then*

$$|S_z| \geq q^f \prod_{j \in \mathcal{J}} \left(1 + q(p-1) 1_{\{z_j \neq 0\}} \right).$$

Consequently the number A_ℓ^{sym} of ordered nontrivial symmetric ℓ -term progressions satisfies

$$A_\ell^{\text{sym}} \geq q^n \left[q^f \prod_{j \in \mathcal{J}} \left(1 + q(p-1)(1 - q^{-t_j}) \right) - 1 \right].$$

In particular, the presence of one such block of degree t forces an excess of at least

$$q^{2n-B-t}(q^t - 1)q(p-1)$$

above the value $q^n(q^f - 1)$ supplied solely by ignored coordinates.

Proof. At a fixed center, set the direction in every unselected controlled block to zero. In a selected block j , always allow direction zero; when $z_j \neq 0$, additionally allow the $q(p-1)$ directions given by the preceding proposition. Each selected block then has a constant norm along every prime-field parameter, so their sums in each output coordinate are constant too. The centered parameters for the standard even ℓ -term progression belong to $\mathbb{F}_p$. Every chosen direction therefore satisfies all the required reflection equalities. The ignored coordinates admit arbitrary directions, giving q^f choices. All these choices are independent and distinct. This proves the local product bound.

The starts and centered parametrizations are in bijection. Sum $|S_z| - 1$ over centers. For each selected block, z_j is independently uniform in its q^{t_j} -element space; its nonzero probability is $1 - q^{-t_j}$. The averaged product gives the displayed global bound. Taking one block and using $f = n - B$ gives the final expression. $\square$

The first high-degree four-term example. Take $q = p = 5$, $\ell = 4$, and $Q = N_3 + N_5$ on $\mathbb{F}_5^3 \oplus \mathbb{F}_5^5$. The highest-two-odd-degree construction attains $S_0 = \{0\}$: the two reflection equations at $\pm 1/2$ and $\pm 3/2$ form an invertible matrix on the cubic and quintic norm values. Nevertheless the degree-five block forces at least

$$5^3(5^5 - 1)5(5 - 1) = 7\,810\,000$$

ordered nontrivial symmetric four-term progressions. Indeed, at every center whose quintic block is nonzero, it supplies at least twenty nonzero directions with the cubic block frozen. There are $5^3(5^5 - 1)$ such centers. This distinguishes sharpness at one prescribed center from simultaneous global sharpness of a particular construction. It does not preclude a different degree-five polynomial from avoiding these patterns.

III.66 Sources 29 and 28: the odd-order constants c_d^{odd} [write]

Added in the sixth write, 6 October 2026.

The next four sections are about $c_d^{\text{odd}} = \sup \|f\|_{U^2}^4 / \|f\|_{U^d}^4$ over real centered f on groups of odd order, source 03's Research question 1 and source 04's Question 3, settled for $d = 3$ by source 04's Theorem III.3.4 ($c_3^{\text{odd}} = 2^{-1/2}$) and open for $d \geq 4$. *Logically the chapter follows source 04's Section III.3*; it is printed here because no number printed earlier may move. Sections III.67–III.69 are source 29's, the main source; Section III.70 is source 28's, written from the same pin; neither cites the other.

What is proved, in one paragraph. For real f with no Fourier mass on self-inverse characters, $\|f\|_{U^d}^{2^d} \geq 2^{-m} \binom{2^m}{m} \|f\|_{U^2}^{2^d}$, $m = 2^{d-2}$ (source 29's Theorem III.68.1; the moment inequality is sharp,

the norm comparison is not claimed sharp), so $c_4^{\text{odd}} \leq (8/35)^{1/4} = 0.6914416$ and $c_d^{\text{odd}} \leq B_d \downarrow 1/2$ (Corollary III.69.1); an interval model gives $c_d^{\text{odd}} \geq 1/3$ for every d , so the limit lies in $[1/3, 1/2]$ (Proposition III.69.2); with self-inverse mass the profile $K_m(\tau)$ (Theorem III.68.3). Source 28 proves the weaker universal $c_d^{\text{odd}} \leq \sqrt{160/323}$ for $d \geq 4$ (Theorem III.70.1), the lower bound $c_4^{\text{odd}} \geq 0.5587294$ by an exact two-harmonic witness (Corollary III.70.8), and that a single cosine is not a local extremizer for any $d \geq 4$ (Theorem III.70.9). The intake's exact certificate on $\mathbb{Z}/7\mathbb{Z}$ gives $c_4^{\text{odd}} \geq 0.5587725$ (note after Corollary III.70.8). The bracket printed here is therefore $0.5587725 \leq c_4^{\text{odd}} \leq 0.6914416$, among sources 01–33; sources 34, 35 and 45, placed and not yet written, record upper coefficients $81/16$, 5 and $4795/832$, that is $c_4^{\text{odd}} \leq 2/3$ and ≤ 0.6454070 (note after Corollary III.69.1). In the quartic cube term the coefficient $100/\sqrt{2} = 70.71$ at $d = 4$ becomes $100(8/35)^{1/4} = 69.14$ (source 29's (III.69.7); source 28's corollary has 70.38).

Relation to the formal statements. `lemma_3_10` (Sections01_03.lean, line 340) has the exact companion `lemma_3_10_holds` (Proofs03CubeUpper.lean, line 82, at HEAD 5d4718772); no declaration concerns c_d^{odd} . Nothing in this chapter is formalized.

[write] **6 October 2026; sources 34–39.** *Sources 34 and 35 continue this chapter.* Source 34's Section 10 (Section III.71) and source 35's Section 6 (Section III.72) follow source 28's Section 8 at the end of this chapter, so that the chapter stays together. Source 34 proves $\|f\|_{U^2}^4 \leq (2/3) \|f\|_{U^4}^4$, coupling the dominant-pair moment profile of source 29 with the exact dual function of a cosine (source 28's identity, uncited there) and two resonant autocorrelation families; source 35, written the same day from the same pin, proves the weaker coefficient 5 by retaining repeated-entry quadruples in the U^3 Fourier energy, an energy inequality of independent interest. The bracket among sources 01–39 is $0.5587725 \leq c_4^{\text{odd}} \leq 2/3$; source 45, placed and not yet written, records $c_4^{\text{odd}} \leq 0.6454070$. Their introductions are Sections F.39 and F.40; source 34's torsion chapter follows (Section III.73).

[write] **7 October 2026; sources 40–46.** *Source 45 continues this chapter* at the end of Part III (Section III.89), because this chapter is no longer the last of Part III: $4795/832$, so $0.5587725 \leq c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$ among sources 01–46.

III.67 Notation and elementary Fourier identities

Source 29, delivered Section 2.

Let G be a finite abelian group. Averages over G are normalized: $\mathbb{E}_x F(x) = |G|^{-1} \sum_{x \in G} F(x)$. Sums over its character group $\widehat{G}$ use counting measure. We write characters multiplicatively and set

$$\widehat{f}(\chi) = \mathbb{E}_x f(x) \overline{\chi(x)}, \quad \Delta_h f(x) = f(x) \overline{f(x-h)}.$$

For $d \geq 1$, define

$$\|f\|_{U^d}^{2^d} = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{\omega \in \{0,1\}^d} \mathcal{C}^{|\omega|} f(x - \omega \cdot h), \quad (\text{III.67.1})$$

where $\mathcal{C}$ is complex conjugation. In particular, $\|f\|_{U^1} = |\mathbb{E} f|$. The sign of the increments is immaterial under uniform averaging. For $d \geq 2$, the recursive identity, together with the Fourier identity for U^2 , is

$$\|f\|_{U^d}^{2^d} = \mathbb{E}_h \|\Delta_h f\|_{U^{d-1}}^{2^{d-1}}, \quad (\text{III.67.2})$$

$$\|f\|_{U^2}^4 = \sum_{\chi \in \widehat{G}} |\widehat{f}(\chi)|^4. \quad (\text{III.67.3})$$

Repeated Cauchy–Schwarz gives the usual monotonicity $\|f\|_{U^1} \leq \|f\|_{U^2} \leq \dots$ and the multilinear Gowers–Cauchy–Schwarz inequality. These elementary norm properties are also developed in Gowers’s Section 3 [1].

For a real function f , define its autocorrelation and Fourier masses by

$$c_f(h) = \mathbb{E}_x f(x)f(x-h) = \sum_{\chi} \lambda_{\chi} \chi(h), \quad \lambda_{\chi} = |\widehat{f}(\chi)|^2. \quad (\text{III.67.4})$$

Then $\lambda_{\chi} \geq 0$ and $\lambda_{\chi^{-1}} = \lambda_{\chi}$. In particular, although c_f need not be pointwise nonnegative, its Fourier coefficients are nonnegative. We use

$$S = \sum_{\chi} \lambda_{\chi}^2 = \|f\|_{U^2}^4.$$

When $f \neq 0$, $S > 0$. The subgroup of self-inverse characters is $\widehat{G}[2] = \{\chi : \chi^2 = 1\}$. On an odd-order group it consists only of the trivial character, so a centered function has zero Fourier mass on $\widehat{G}[2]$.

III.68 Positive autocorrelation moments and higher norm bounds

Source 29, delivered Section 3.

Theorem III.68.1 (Central-binomial autocorrelation bound). *Let $f : G \rightarrow \mathbb{R}$, and suppose $\widehat{f}(\chi) = 0$ whenever $\chi^2 = 1$. For every integer $m \geq 1$,*

$$\mathbb{E}_h c_f(h)^{2m} \geq \binom{2m}{m} \left(\frac{S}{2}\right)^m. \quad (\text{III.68.1})$$

Consequently, for every $d \geq 2$, with $m = 2^{d-2}$,

$$\|f\|_{U^d}^{2^d} \geq C_m \|f\|_{U^2}^{2^d}, \quad C_m = 2^{-m} \binom{2m}{m}. \quad (\text{III.68.2})$$

The constant in (III.68.1) is optimal uniformly over finite abelian groups. Sharpness is not asserted for (III.68.2) when $d \geq 3$.

Proof. Choose one character χ_j from each nonzero inverse pair, and write $b_j = \lambda_{\chi_j}$. Thus $S = 2 \sum_j b_j^2$. The zero function is immediate, so assume at least one pair occurs. Expanding the Fourier series of c_f^{2m} and averaging retains precisely the monomials whose product character is trivial. Their coefficients are nonnegative. We may therefore retain only monomials that use χ_j and χ_j^{-1} equally often for every j . If each occurs k_j times, where $\sum_j k_j = m$, the contribution is

$$\frac{(2m)!}{\prod_j (k_j!)^2} \prod_j b_j^{2k_j}.$$

Hence

$$\begin{aligned}
\mathbb{E}c_f^{2m} &\geq \binom{2m}{m} \sum_{\sum k_j = m} \binom{m}{k_1, k_2, \dots}^2 \prod_j b_j^{2k_j} \\
&\geq \binom{2m}{m} \sum_{\sum k_j = m} \binom{m}{k_1, k_2, \dots} \prod_j b_j^{2k_j} \\
&= \binom{2m}{m} \left(\sum_j b_j^2 \right)^m.
\end{aligned} \tag{III.68.3}$$

The second inequality only uses the fact that each multinomial coefficient is an integer at least one.

By (III.67.2) and $\|g\|_{U^{d-1}} \geq |\mathbb{E}g|$,

$$\|f\|_{U^d}^{2^d} \geq \mathbb{E}_h |c_f(h)|^{2^{d-1}}.$$

Here $2^{d-1} = 2m$, and $S^m = \|f\|_{U^2}^{2^d}$, proving the norm comparison. To see moment sharpness, take a character of order greater than $2m$ and $f = 2 \operatorname{Re}(a\chi)$, with $a \neq 0$. The only surviving monomials in c_f^{2m} then contain the two characters equally often, and equality holds. $\square$

III.68.1 Equality and stability at the moment level

The moment argument contains more information than its last scalar bound. It quantifies the cost of distributing Fourier fourth-power mass among several inverse pairs.

Proposition III.68.2 (Pair concentration). *Under the assumptions of Theorem III.68.1, let $f \neq 0$, $m \geq 2$, and*

$$p_j = \frac{2\lambda_{\chi_j}^2}{S}, \quad \sum_j p_j = 1.$$

Then

$$\frac{\mathbb{E}c_f^{2m}}{C_m S^m} \geq m - (m-1) \sum_j p_j^m. \tag{III.68.4}$$

If the left side is at most $1 + \varepsilon$, where $0 \leq \varepsilon < m-1$, some inverse pair has mass

$$p_j \geq \left(1 - \frac{\varepsilon}{m-1}\right)^{1/(m-1)}. \tag{III.68.5}$$

For that pair, putting $a = \widehat{f}(\chi_j)$,

$$\|f - 2 \operatorname{Re}(a\chi_j)\|_{U^2}^4 \leq S \left[1 - \left(1 - \frac{\varepsilon}{m-1}\right)^{1/(m-1)}\right]. \tag{III.68.6}$$

Equality in (III.68.1), for $m \geq 2$, holds exactly when the Fourier support is one inverse pair and, for its character order n , $n/\gcd(n, 2) > m$.

Proof. For a composition k of m , the multinomial coefficient is one if all mass is in a single coordinate, and is at least m otherwise. The normalized balanced sum in (III.68.3) is the expectation of this

multinomial coefficient under the multinomial distribution with probabilities p_j . The probability of a pure composition is $\sum_j p_j^m$. This proves (III.68.4). It gives $\sum_j p_j^m \geq 1 - \varepsilon/(m-1)$, whereas $\sum_j p_j^m \leq (\max_j p_j)^{m-1}$. Removing the largest pair and using (III.67.3) proves the distance statement.

At equality every mixed composition must vanish, so just one pair remains. For that pair the extra monomials have exponent $2j - 2m$, $0 \leq j \leq 2m$, on the chosen character. A nonzero such exponent can be divisible by n exactly when $n/\gcd(n, 2) \leq m$. Its coefficient is positive. This proves the stated criterion. $\square$

This is stability for the autocorrelation moment and approximation in U^2 . It neither asserts approximation in L^2 nor classifies near-minimizers of $\|f\|_{U^d}/\|f\|_{U^2}$. The first inequality in the norm proof can have its own nonzero gap.

III.68.2 An explicit profile for self-inverse Fourier mass

The absence of order-two characters is essential for a uniform improvement over one: a nontrivial real character has all Gowers norms equal to one. Its contribution can nevertheless be tracked exactly in the moment method.

Theorem III.68.3 (Torsion-dependent moment profile). *Let $f : G \rightarrow \mathbb{R}$ be nonzero, with no restriction on its mean or on G . Set*

$$\tau = \frac{\sum_{\chi^2=1} \lambda_\chi^2}{S} \in [0, 1],$$

and define

$$K_m(\tau) = \sum_{k=0}^m \binom{2m}{2k} \binom{2k}{k} \tau^{m-k} \left(\frac{1-\tau}{2}\right)^k. \quad (\text{III.68.7})$$

For every integer $m \geq 1$,

$$\mathbb{E}_h c_f(h)^{2m} \geq K_m(\tau) S^m. \quad (\text{III.68.8})$$

For $m = 2^{d-2}$, it follows that $\|f\|_{U^d}^{2d} \geq K_m(\tau) \|f\|_{U^2}^{2d}$. For each m and each τ , the moment constant in (III.68.8) is attained on some finite abelian group.

Proof. Separate the self-inverse frequencies from the remaining inverse pairs. In the Fourier expansion retain terms in which every self-inverse frequency occurs an even number of times and each non-self-inverse pair occurs in equal numbers. These terms are exactly the moment of

$$X + Y = \sum_i a_i \epsilon_i + \sum_j 2b_j \cos \theta_j,$$

where the ϵ_i are independent uniform signs and the θ_j are independent uniform circle variables, all mutually independent. Their variances are $\mathbb{E}X^2 = \tau S$ and $\mathbb{E}Y^2 = (1-\tau)S$. All other actual Fourier terms are nonnegative, so discarding them is legitimate even when the original characters have extra relations.

Odd proxy moments vanish. Jensen's inequality gives $\mathbb{E}X^{2r} \geq (\tau S)^r$, including $r = 0$. The balanced multinomial calculation gives $\mathbb{E}Y^{2k} \geq \binom{2k}{k} ((1-\tau)S/2)^k$. Expanding $\mathbb{E}(X+Y)^{2m}$ and applying these two estimates gives (III.68.7).

For sharpness, use $G = C_2 \times C_p$, with p an odd prime greater than $2m$. Put Fourier squared mass $\lambda_\sigma = \sqrt{\tau S}$ on the nontrivial character of C_2 , and $\lambda_\chi = \lambda_{\chi^{-1}} = \sqrt{(1-\tau)S/2}$ on one pair from

C_p . A real function with these squared Fourier coefficients exists by taking nonnegative square roots as Fourier coefficients. Only the retained moments survive, and both component estimates are equalities. $\square$

The endpoints are $K_m(0) = C_m$ and $K_m(1) = 1$. The profile need not be monotone between them: for example, $K_2(\tau) = \frac{3}{2} + 3\tau - \frac{7}{2}\tau^2$. It describes how Fourier mass in both sectors can increase an even moment; it is not a monotone penalty for torsion alone.

III.69 Consequences for the unresolved higher norm constants

Source 29, delivered Section 4.

For $d \geq 3$, define

$$c_d^{\text{odd}} = \sup_{\substack{|G| \text{ odd}, f: G \rightarrow \mathbb{R} \\ \mathbb{E}f=0, f \neq 0}} \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4}. \quad (\text{III.69.1})$$

The supremum ranges over finite groups as well as functions. The existing report proves $c_3^{\text{odd}} = 2^{-1/2}$ [135]. Monotonicity alone therefore gives that same upper bound for every $d \geq 3$. In norm-power form its coefficient is $2^{m/2}$, where $m = 2^{d-2}$.

Corollary III.69.1 (A strict improvement from dimension four onward). *For $d \geq 4$ and $m = 2^{d-2}$,*

$$c_d^{\text{odd}} \leq B_d := \frac{2}{\binom{2m}{m}^{1/m}} < 2^{-1/2}. \quad (\text{III.69.2})$$

In particular,

$$\|f\|_{U^4}^{16} \geq \frac{35}{8} \|f\|_{U^2}^{16} \quad (\text{III.69.3})$$

for real centered f on odd-order groups. The bounds B_d decrease to $1/2$.

[write] **6 October 2026; sources 28–33.** *The best upper coefficient among sources 01–33, not the best known.* Among sources 01–33, $35/8$ is the best coefficient for real centered f on odd-order groups at $d = 4$ (source 28’s universal bound gives $(323/160)^2 = 4.0754\dots$, note after its Theorem III.70.1). Sources placed in this directory and not yet written in record larger coefficients: source 35 records 5, source 34 records $81/16$, that is $c_4^{\text{odd}} \leq 2/3$, and source 45 records $4795/832$, that is $c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$ (their placements, checked at intake; for the writes that add them). With source 28’s lower bound and the intake’s certificate the bracket printed by this write is $0.5587725 \leq c_4^{\text{odd}} \leq (8/35)^{1/4} = 0.6914416$ (item III.1). The decimals of the table were recomputed at the write.

[write] **6 October 2026; sources 34–39.** Sources 34 and 35 are now written in: source 34’s Theorem III.71.1 ($81/16$, $c_4^{\text{odd}} \leq 2/3$, the best among sources 01–39) and source 35’s Theorem III.72.1 (5) follow source 28’s Section 8 at the end of this chapter; source 45 ($4795/832$) is placed and not yet written. For $5 \leq d \leq 8$ source 34’s Remark III.71.10 improves the bounds B_d of the table above slightly (recomputed at intake); for $d \geq 9$ they stand.

[write] **7 October 2026; sources 40–46.** Source 45 is now written in: its Theorem III.90.1 ($4795/832$, $c_4^{\text{odd}} \leq 0.6454070$, the best among sources 01–46) is printed at the end of Part III (Section III.89), and source 49, placed by `f8bc5e2ec` and not yet written in, records $29/4$, that is $c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ (for its write); by monotonicity it also bounds c_d^{odd} for $d \geq 5$, but source 34’s β_d for $5 \leq d \leq 8$ and the B_d of the table for $d \geq 9$ are smaller.

[write] **7 October 2026; sources 47–53.** Source 49 is now written in (Section III.92): $c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$, the best among sources 01–53, and $c_5^{\text{odd}} \leq 29/50 = 0.58$; for $d \geq 6$ source 34's β_d ($5 \leq d \leq 8$) and the B_d of the table ($d \geq 9$) are smaller.

Proof. The norm theorem gives $c_d^{\text{odd}} \leq C_m^{-1/m} = B_d$. At $m = 4$, $C_4 = 35/8 > 4 = 2^{4/2}$. Vandermonde's identity includes the central summand $\binom{4m}{2m} \geq \binom{2m}{m}^2$; hence $C_{2m} \geq C_m^2$. Induction along powers of two proves the strict comparison with $2^{m/2}$, and also the monotonicity of B_d . Finally $4^m/(2m+1) \leq \binom{2m}{m} \leq 4^m$, which gives the limit. $\square$

d	m	Moment coefficient C_m	Upper bound for c_d^{odd}
3	2	3/2 (use the older sharp coefficient 2)	0.707106781
4	4	35/8	0.691441569
5	8	6435/128	0.612819643
6	16	300540195/32768	0.565388162
7	32	$2^{-32} \binom{64}{32}$	0.537413970
8	64	$2^{-64} \binom{128}{64}$	0.521168336

Table III.69.1: The decimals are approximations to the proved exact upper bounds. The new moment estimate is weaker at $d = 3$, where the previously proved sharp result is retained.

III.69.1 A lower bound uniform in the dimension

Proposition III.69.2 (A uniform interval model). *For every $d \geq 3$, $c_d^{\text{odd}} \geq 1/3$. Consequently the nonincreasing sequence c_d^{odd} has a limit in $[1/3, 1/2]$.*

[write] **6 October 2026; sources 28–33.** This answers the part of source 28's Research question on the asymptotic constants that asks for a positive lower bound uniform in d (item III.1); checked at the write: the ratio $n(n^2 + n + 1)/(3(n + 1)^3)$ is 0.278, 0.327, 0.3327 at $n = 10, 100, 1000$.

[write] **7 October 2026; sources 60–69.** *The limit is $1/3$ (source 63, Section III.94): $c_d^{\text{odd}} \leq 1/3 + 10 \cdot 2^{-\lfloor d/8 \rfloor}$ for $d \geq 80$ (its Theorem III.97.1), from the sharp bounded inequality $\|f\|_{U^2}^4 \leq \|f\|_{\infty}^4/3$ on odd-order groups; source 60 had shown that the limit is at most 0.3464533951.*

Proof. On $G = \mathbb{Z}/(2n+1)\mathbb{Z}$, let $A = \{0, \dots, n-1\}$, $\delta = n/(2n+1)$, and $f = 2(\mathbf{1}_A - \delta)$. There is no wraparound in sums of two members of A . The additive energy is

$$E(A) = n^2 + 2 \sum_{j=1}^{n-1} j^2 = \frac{2n^3 + n}{3}.$$

By Fourier orthogonality,

$$\|f\|_{U^2}^4 = 16 \left(\frac{E(A)}{(2n+1)^3} - \delta^4 \right), \quad \|f\|_{\infty} = \frac{2(n+1)}{2n+1}.$$

Since $\|f\|_{U^d} \leq \|f\|_{\infty}$,

$$\frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4} \geq \frac{n(n^2 + n + 1)}{3(n+1)^3} \rightarrow \frac{1}{3}.$$

This works for each fixed d , with a bound independent of d . Monotonicity and Corollary III.69.1 prove the limit assertion. $\square$

III.69.2 A more informative lower bound in dimension four

For a prime $p > 16$ and $f = a\chi + \bar{a}\chi^{-1}$ on $\mathbb{Z}/p\mathbb{Z}$, with χ nontrivial, one has

$$\|f\|_{U^4}^{16} = 222|a|^{16}, \quad \|f\|_{U^2}^{16} = 16|a|^{16}. \quad (\text{III.69.4})$$

Here is a finite combinatorial verification of the coefficient. A surviving Fourier monomial chooses eight positive-sign vertices of $\{0, 1\}^4$, with four ones in each coordinate. Split by the fourth coordinate into four-subsets of the three-cube. If $n(v)$ counts four-subsets with coordinate-sum vector v , complementing a subset shows $n(4\mathbf{1} - v) = n(v)$, and the required number is $\sum_v n(v)^2$. The following complete classification uses coordinate permutations and coordinate reflections:

sorted $(v_1 - 2 , v_2 - 2 , v_3 - 2)$	number of vectors v	$n(v)$
(0, 0, 0)	1	8
(0, 0, 1)	6	4
(0, 0, 2)	6	1
(0, 1, 1)	12	2
(1, 1, 1)	8	1

For example, a coordinate sum zero or four determines a full face; for sums $(1, 2, 2)$, the unique point on one side of the first coordinate determines the omitted point on the other side, giving four choices. For $(1, 1, 2)$, separating whether the first two coordinate ones occur at the same point gives two choices. For $(1, 1, 1)$ the subset is forced. The balanced case $(2, 2, 2)$ consists of six unions of two antipodal pairs and two parity tetrahedra. Thus the sum is $64 + 96 + 6 + 48 + 8 = 222$. The condition $p > 16$ prevents nonzero integer exponents from vanishing modulo p .

It follows that

$$\left(\frac{8}{111}\right)^{1/4} \leq c_4^{\text{odd}} \leq \left(\frac{8}{35}\right)^{1/4}. \quad (\text{III.69.5})$$

The endpoints are approximately 0.518133595 and 0.691441569. The interval has not been closed. In particular, the number $35/8$ is a rigorously improved coefficient, not the asserted sharp answer.

[write] **6 October 2026; sources 28–33.** The lower end of (III.69.5) is the single-cosine value. Source 28, from the same pin, raises it to 0.5587294 by an exact two-harmonic witness (Corollary III.70.8), and the intake's certificate on $\mathbb{Z}/7\mathbb{Z}$ to 0.5587725 (note there); source 28's Theorem III.70.9 shows that a single inverse pair is not extremal for any $d \geq 4$.

III.69.3 How the coefficient enters centered cube counting

For $Q_d(F) = \|F\|_{U^d}^{2d}$ and real centered f , expansion of $Q_d(\delta + f)$ has no terms with one, two, or three selected vertices: at those vertices the affine forms are independent. On odd-order groups the fourth-degree term is

$$P_d \delta^{2^d-4} \|f\|_{U^2}^4, \quad P_d = 2^{d-3}(3^d - 2^{d+1} + 1). \quad (\text{III.69.6})$$

This coefficient and the torsion qualifications are proved in the existing cube-expansion reports [135, 137]. The new norm estimate therefore changes the available bound on this term to

$$P_d B_d \delta^{2^d-4} \|f\|_{U^d}^4, \quad d \geq 4. \quad (\text{III.69.7})$$

For $d = 4$, $P_4 = 100$: the coefficient improves from $100/\sqrt{2} \approx 70.710678$ to $100(8/35)^{1/4} \approx 69.144157$. This is an explicit improvement to one local coefficient. Terms of higher degree and subsequent inverse and localization losses still have to be controlled before a final progression-free-set bound could change.




Figure III.69.1: Left: the proved upper bounds for c_d^{odd} , with the dimension-independent lower bound $1/3$ and the sharper cosine lower bound at $d = 4$. Right: the full fourth-dimensional moment profile, expressed as $K_4(\tau)^{-1/4}$. The curve is a valid bound from the moment method, not a claimed sharp Gowers-norm constant.

III.70 A stronger real norm gap and a higher-order obstruction

Source 28, delivered Section 8.

Delivered part heading of source 28: **Strict norm gaps and the failure of cosine extremality**

The earlier repository report proves the sharp real, mean-zero comparison $\|f\|_{U^3}^8 \geq 2\|f\|_{U^2}^8$ on odd-order groups, with equality precisely for a single conjugate pair of nonzero Fourier coefficients [132]. Its higher-order analogue was left open. We improve the available universal bound at every order at least four, and show that the single-character equality model does not extend to these orders. The two conclusions are independent: the universal estimate is a lower bound for a norm, whereas the explicit examples give lower bounds for the extremal ratio.

All expectations in this section use the uniform probability measure. For a real function on a finite abelian group, our convention is

$$\|f\|_{U^d}^{2d} = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{\omega \in \{0,1\}^d} f(x + \omega \cdot h). \quad (\text{III.70.1})$$

The usual conjugations in the complex definition have no effect here. Fourier coefficients are normalized by $\hat{f}(\xi) = \mathbb{E}_x f(x) \overline{\xi(x)}$. In particular, $\|f\|_{U^2}^4 = \sum_{\xi} |\hat{f}(\xi)|^4$. Write

$$c_d^{\text{odd}} = \sup_{\substack{G \text{ finite abelian of odd order} \\ f: G \rightarrow \mathbb{R}, \mathbb{E}f=0, f \neq 0}} \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4}. \quad (\text{III.70.2})$$

No boundedness restriction changes this supremum, since the ratio is invariant under multiplying f by a nonzero scalar.

Theorem III.70.1 (Explicit improvement beyond the third norm). *For every finite abelian group of odd order, every real mean-zero function f , and every integer $d \geq 4$,*

$$\|f\|_{U^d}^4 \geq \sqrt{\frac{323}{160}} \|f\|_{U^2}^4. \quad (\text{III.70.3})$$

Moreover,

$$\left(\frac{133153321267264}{1366299579065739} \right)^{1/4} \leq c_4^{\text{odd}} \leq \sqrt{\frac{160}{323}}. \quad (\text{III.70.4})$$

The endpoints are approximately 0.558729390985 and 0.703815343078, respectively. The upper endpoint strictly improves the previously available $2^{-1/2} = 0.707106781187\dots$

[write] **6 October 2026; sources 28–33.** *Superseded the same day by source 29.* The universal bound (III.70.3) is $\|f\|_{U^4}^{16} \geq (323/160)^2 \|f\|_{U^2}^{16} = 4.0754\dots \|f\|_{U^2}^{16}$ at $d = 4$; source 29’s Corollary III.69.1, written from the same pin and citing neither this source nor being cited by it, gives $\|f\|_{U^4}^{16} \geq (35/8) \|f\|_{U^2}^{16}$, hence $c_4^{\text{odd}} \leq (8/35)^{1/4} = 0.6914416$ against $\sqrt{160/323} = 0.7038153$ here, and $c_d^{\text{odd}} \leq B_d < 0.6914416$ for every $d \geq 5$. The sentence that the upper endpoint “strictly improves the previously available $2^{-1/2}$ ” was true at the pin. The lower endpoint 0.5587294 is this source’s own; it is improved slightly by the intake’s certificate (note after Corollary III.70.8). The bracket is item III.1; later placed sources record further upper bounds (note after source 29’s Corollary III.69.1).

We give all ingredients below. The numerical displays are descriptive; the assertions use the exact rational quantities in the formulas. Neither endpoint is asserted to be sharp.

III.70.1 A short bounded example on five points

Proposition III.70.2 (A rational counterexample to the cosine value). *On $\mathbb{Z}/5\mathbb{Z}$, let*

$$(f(0), f(1), f(2), f(3), f(4)) = (0, 1, 1, -1, -1).$$

Then $\mathbb{E}f = 0$, $|f| \leq 1$, and

$$\|f\|_{U^2}^4 = \frac{36}{125}, \quad \|f\|_{U^4}^{16} = \frac{52}{625}, \quad \frac{\|f\|_{U^4}^{16}}{\|f\|_{U^2}^{16}} = \frac{5078125}{419904} < \frac{111}{8}. \quad (\text{III.70.5})$$

A single real character pair of odd order at least five has the ratio $111/8$. Thus it cannot describe the general fourth-order optimum.

Proof. The unnormalized autocorrelations of f are $(4, 1, -3, -3, 1)$. Their squared sum is 36, giving the first identity. For $h, k \in \mathbb{Z}/5\mathbb{Z}$, put

$$g_{h,k}(x) = f(x)f(x+h)f(x+k)f(x+h+k).$$

Twice applying the derivative identity for Gowers norms gives $\|f\|_{U^4}^{16} = \mathbb{E}_{h,k} \|g_{h,k}\|_{U^2}^4$. For an integer-valued g , write

$$E(g) = \sum_t \left(\sum_x g(x)g(x+t) \right)^2, \quad \|g\|_{U^2}^4 = E(g)/125.$$

The possibilities are

Condition on (h, k)	Number of pairs	$E(g_{h,k})$
$h = k = 0$	1	52
Exactly one of h, k is zero	8	19
$h, k \neq 0$ and $h = \pm k$	8	6
$h, k \neq 0$ and $h \neq \pm k$	8	1

Indeed, the first function is the indicator of a four-point set. In the second case it is the indicator of a three-point set. Its nonzero-shift intersection counts are 2, 2, 1, 1, so its energy is $3^2 + 2 \cdot 2^2 + 2 \cdot 1^2 = 19$. In the last two cases its support has size two and one, respectively, with values in $\{-1, 1\}$; the corresponding energies are 6 and 1, independently of the signs. The sum of the energies is $52 + 8 \cdot 19 + 8 \cdot 6 + 8 = 260$. Dividing by 5^5 gives $52/625$, and the ratio follows. $\square$

III.70.2 The exact fourth-order dual function of a cosine

For a real function g , define

$$D_4g(x) = \mathbb{E}_{h_1, h_2, h_3, h_4} \prod_{\substack{\omega \in \{0,1\}^4 \\ \omega \neq 0}} g(x + \omega \cdot h). \quad (\text{III.70.6})$$

The Gowers–Cauchy–Schwarz inequality implies

$$|\mathbb{E}_x f(x) D_4g(x)| \leq \|f\|_{U^4} \|g\|_{U^4}^{15}. \quad (\text{III.70.7})$$

This finite multilinear inequality, rather than an unproved description of a dual norm, is all that is needed.

Lemma III.70.3 (Exact cosine counts). *Let χ be a character of odd order at least five, let $\theta \in \mathbb{R}$, and write $g_j = e^{ij\theta} \chi^j + e^{-ij\theta} \chi^{-j}$. Then*

$$\|g_1\|_{U^4}^{16} = 222, \quad D_4g_1 = 111g_1 + 27g_3 + g_5. \quad (\text{III.70.8})$$

The formula remains valid when χ has order five; in that case g_5 is constant.

Proof. Expand each factor in characters. In the average defining D_4g_1 , the coefficient choices are signs on the fifteen nonzero Boolean vertices. Each direction average requires the sum of eight signs to be divisible by the order of χ . Such a sum is even and has absolute value at most eight. For an odd order at least five, it must therefore be zero as an integer. The same observation applies to both halves of a full four-dimensional cube; its total frequency is then also zero.

First count the full-cube choices. Split the cube into two three-dimensional faces. Each face has four plus signs. If $n(t)$ is the number of four-subsets of $\{0, 1\}^3$ whose coordinate sum is $t = (t_1, t_2, t_3)$, the full count is $\sum_t n(t)^2$. The required finite generating identity is

$$[z^4] \prod_{\epsilon \in \{0,1\}^3} (1 + zx^\epsilon) = x_1 x_2 x_3 \prod_{j=1}^3 (1 + x_j)^2 + \sum_{i=1}^3 (1 + x_i^4) \prod_{j \neq i} x_j^2. \quad (\text{III.70.9})$$

Here $x^\epsilon = \prod_j x_j^{\epsilon_j}$. This identity follows by multiplying the eight binomials and collecting the coefficient of z^4 . It gives the following complete table:

Type of t	Number of such t	$n(t)$
One coordinate 0 or 4, the other two 2	6	1
All three coordinates in $\{1, 3\}$	8	1
Two in $\{1, 3\}$, one equal to 2	12	2
One in $\{1, 3\}$, two equal to 2	6	4
$(2, 2, 2)$	1	8

Thus the count is $6 + 8 + 12 \cdot 4 + 6 \cdot 16 + 64 = 222$. The factor $e^{i\theta \sum_{\omega} s_{\omega}}$ is one for every such choice, so this proves the first identity in (III.70.8).

For the dual function, a choice of signs on the fifteen nonzero vertices contributes frequency $\sum s_{\omega}$. The counts at frequencies 1 and -1 are both 111: adjoining the missing sign at the origin produces the 222 full-cube assignments, divided equally by its sign. The counts at frequencies 3 and -3 are both 27, as follows. Place a frequency 3 at the missing origin, so that the other fifteen signs must have total -3 . Their six plus-sign vertices form a set P with four selected vertices on each coordinate face. After complementing all four coordinates, we obtain six distinct vertices, excluding $(1, 1, 1, 1)$, with coordinate sums all equal to two. Let a_j count their vertices of Hamming weight j . Then

$$a_0 + a_1 + a_2 + a_3 = 6, \quad a_1 + 2a_2 + 3a_3 = 8, \quad a_0 \in \{0, 1\}, \quad a_1 \leq 4.$$

Exactly three possibilities remain:

a_0	a_1	a_2	a_3	Number
0	4	2	0	3
1	2	3	0	12
1	3	1	1	12

In the first case the two edges form a perfect matching on four coordinates. In the second, choose the two singleton coordinates in six ways; the three edges form a path with those endpoints, in two ways. In the third, choose the triple in four ways and choose which of its three coordinates has no singleton; the edge is then forced. This proves $27 = 3 + 12 + 12$.

At frequency -5 , the five plus-sign vertices must have a total of sixteen coordinate incidences. The only possibility is the vertex of weight four together with the four vertices of weight three. Hence the count is one; the frequency 5 count is also one. Frequencies of absolute value at least seven are impossible, since four distinct vertices have at most $4 + 3 + 3 + 3 = 13$ incidences. These counts exhaust all odd frequencies, and prove the dual formula. $\square$

[write] **6 October 2026; sources 34–39.** Also proved by source 34 (Lemma III.71.3) and recounted by source 35 (Lemma III.72.3), neither citing this source.

Lemma III.70.4 (The order-three case). *If χ has order three, then for $g = e^{i\theta}\chi + e^{-i\theta}\chi^{-1}$,*

$$\begin{aligned} \|g\|_{U^4}^{16} &= 286 + 32 \cos(6\theta) + 16 \cos(12\theta) \\ &= 262 + 32 \left(\cos(6\theta) + \frac{1}{2} \right)^2 \geq 262. \end{aligned} \tag{III.70.10}$$

Proof. Let P be the plus-sign vertices in a character expansion of the four-cube. The congruences require $|P| \equiv 2 \pmod{3}$ and require each coordinate face to contain 1, 4, or 7 vertices of P . The possibilities are

$ P $	Total frequency $2 P - 16$	Number
2	-12	8
5	-6	16
8	0	286
11	6	16
14	12	8

For two vertices, every coordinate must occur once, so the vertices are antipodal, giving eight pairs. For five vertices, reflect every coordinate having four selected vertices. Each coordinate now occurs once, forcing the set consisting of the zero vertex and the four unit vertices. Its sixteen reflected copies are distinct. For eight vertices with four on every face, the count is 222 by (III.70.9). Otherwise some coordinate has one or seven selected vertices. Reflect to make this number one. The set is one entire three-face, with one vertex removed, and the corresponding vertex inserted across that coordinate. The remaining coordinate counts must all be four. There are $4 \cdot 2 \cdot 8 = 64$ such sets, each with exactly one exceptional coordinate. This gives $222 + 64 = 286$. Taking complements supplies the last two rows. Summing the associated phases proves the first line of (III.70.10); the second follows from the double-angle identity. $\square$

III.70.3 An explicit universal improvement

We recall precisely which earlier estimate is used. Let $f \neq 0$ be real and mean-zero on an odd-order group, set

$$S = \sum_{\xi} |\widehat{f}(\xi)|^4 = \|f\|_{U^2}^4, \quad p = \max_{\xi \neq 0} \frac{2|\widehat{f}(\xi)|^4}{S}.$$

The nonzero Fourier coefficients occur in distinct conjugate pairs, so $0 < p \leq 1$.

Lemma III.70.5 (Two bounds retaining the largest pair). *With this notation, put $R = \|f\|_{U^4}^{16}/S^4$. Then*

$$R \geq (5 - 3p)^2. \quad (\text{III.70.11})$$

If a maximizing character has order at least five, also

$$R \geq \frac{111}{8} \left[p^{1/4} - \frac{28}{111}(1-p)^{1/4} \right]_+^{16}, \quad (\text{III.70.12})$$

where $[t]_+ = \max(t, 0)$. If it has order three, then

$$R \geq \frac{131}{8} p^4. \quad (\text{III.70.13})$$

Proof. The prior real U^2 – U^3 argument retains the stronger intermediate inequality

$$\|f\|_{U^3}^8 \geq 5S^2 - 6M, \quad M = \sum_{\xi} |\widehat{f}(\xi)|^8.$$

For completeness, write $\lambda_{\xi} = |\widehat{f}(\xi)|^2$. Writing $a_r = \widehat{f}(r)$, the exact Fourier identity is

$$\|f\|_{U^3}^8 = \sum_{s,t} \left| \sum_r a_r \overline{a_{r+s}} \overline{a_{r+t}} a_{r+s+t} \right|^2.$$

Retaining its two coordinate axes gives $\|f\|_{U^3}^8 \geq 2E(\lambda) - S^2$. In the additive energy $E(\lambda)$, retain the union of $r = t, s = w$; $r = w, s = t$; and $s = -r, w = -t$. Each contributes S^2 , each pairwise intersection contributes M , and their triple intersection has zero weight because the group has odd order and $\hat{f}(0) = 0$. Thus $E(\lambda) \geq 3S^2 - 3M$, proving the displayed estimate. If p_j are the normalized conjugate-pair fourth-power masses, then $2M/S^2 = \sum_j p_j^2 \leq \max_j p_j = p$. Consequently $\|f\|_{U^3}^8 \geq (5 - 3p)S^2$. Norm monotonicity $\|f\|_{U^4} \geq \|f\|_{U^3}$ proves (III.70.11).

Choose a maximizing character χ and write $\hat{f}(\chi) = ae^{i\theta}$, where $a = (pS/2)^{1/4}$. If its order is at least five, apply (III.70.7) to the function g_1 in Lemma III.70.3. The frequencies $\chi^{\pm 3}$ and $\chi^{\pm 5}$ do not coincide with the maximizing pair. A trivial fifth power contributes zero because f is centered. Every remaining coefficient has modulus at most $((1 - p)S/2)^{1/4}$. Therefore

$$\mathbb{E}fD_4g_1 \geq 222(pS/2)^{1/4} - 56((1 - p)S/2)^{1/4}.$$

Since $\|g_1\|_{U^4}^{16} = 222$, inequality (III.70.7), followed by taking the sixteenth power and dividing by S^4 , gives (III.70.12).

If χ has order three, average f over translates by $\ker \chi$. This operation does not increase a Gowers norm, by its triangle inequality and translation invariance. The quotient has three points, and the centered averaged function is exactly $a(e^{i\theta}\chi + e^{-i\theta}\chi^{-1})$. Lemma III.70.4 gives $\|f\|_{U^4}^{16} \geq 262a^{16} = (131/8)p^4S^4$. $\square$

Upper bound in Theorem III.70.1. The zero function is immediate. Otherwise use Lemma III.70.5 and split at $p_0 = 159/160$. If $p \leq p_0$, then

$$R \geq (5 - 3p_0)^2 = (323/160)^2.$$

If $p > p_0$ and the maximizing character has order three,

$$R \geq (131/8)(159/160)^4 > (323/160)^2.$$

In the remaining case use the exact rational comparisons

$$(599/600)^4 < 159/160, \quad 1/160 < (9/32)^4.$$

They imply

$$p^{1/4} - \frac{28}{111}(1 - p)^{1/4} > \frac{599}{600} - \frac{28}{111} \frac{9}{32} = \frac{5147}{5550}.$$

Finally,

$$\frac{111}{8} \left(\frac{5147}{5550} \right)^{16} > \left(\frac{323}{160} \right)^2.$$

For example, the last comparison is the integer inequality $111 \cdot 3200 \cdot 5147^{16} > 104329 \cdot 5550^{16}$. Thus $R \geq (323/160)^2$ in every case, which is (III.70.3) for $d = 4$. Norm monotonicity extends it to all $d \geq 4$. Rearranging proves the upper bound for c_4^{odd} . $\square$

Corollary III.70.6 (A smaller local quartic cube coefficient). *Let $d \geq 4$, $q = 2^d$, $\delta \geq 0$, and let f be real and centered on an odd-order group. Put $u = \|f\|_{U^d}$ and*

$$P_d = 2^{d-3}(3^d - 2^{d+1} + 1).$$

The exact quartic expansion established in the earlier repository report [132] now gives

$$\left| \mathbb{E}_{x,h} \prod_{\omega \in \{0,1\}^d} (\delta + f(x + \omega \cdot h)) - \delta^q \right| \leq P_d \sqrt{\frac{160}{323}} \delta^{q-4} u^4 + \sum_{r=5}^q \binom{q}{r} \delta^{q-r} u^r. \quad (\text{III.70.14})$$

For $d = 4$, the available leading coefficient decreases from $100/\sqrt{2}$ to $100\sqrt{160/323}$, approximately 70.3815343.

Proof. The existing exact quartic term is $P_d \delta^{q-4} \|f\|_{U^2}^4$. Substitute (III.70.3). The higher support terms retain their original Gowers–Cauchy–Schwarz bounds. $\square$

III.70.4 An exact two-harmonic family and a stronger lower bound

Let χ have order greater than 24, and set

$$f_b = \chi + \chi^{-1} + b(\chi^3 + \chi^{-3}), \quad b \in \mathbb{R}. \quad (\text{III.70.15})$$

These functions are real and centered. Orthogonality immediately gives

$$\|f_b\|_{U^2}^4 = 2(1 + b^4). \quad (\text{III.70.16})$$

Proposition III.70.7 (Exact fourth-norm polynomial). *For (III.70.15),*

$$\|f_b\|_{U^4}^{16} = P(b) = \sum_{j=0}^{16} c_j b^j, \quad (\text{III.70.17})$$

where the coefficients are

j	c_j	j	c_j
0	222	9	28320
1	864	10	24160
2	4080	11	4256
3	9120	12	5920
4	25640	13	0
5	31808	14	0
6	55120	15	0
7	49280	16	222
8	52360		

In particular, if $R_4(b) = \|f_b\|_{U^4}^{16} / \|f_b\|_{U^2}^{16}$, then

$$R_4(0) = \frac{111}{8}, \quad R'_4(0) = 54 > 0. \quad (\text{III.70.18})$$

Proof. A Fourier choice attaches $s_\omega \in \{-3, -1, 1, 3\}$ to each four-cube vertex and has weight $b^{\#\{\omega: |s_\omega|=3\}}$. Orthogonality requires

$$\sum_{\omega} s_{\omega} = 0, \quad \sum_{\omega} \omega_i s_{\omega} = 0 \quad (1 \leq i \leq 4),$$

initially as congruences modulo the character order. Each half-cube sum has absolute value at most 24. The total and first-coordinate congruences imply that both first-coordinate half sums are zero integers; the other coordinate sums are also zero integers. Thus the order assumption removes every congruence alias.

Here is a finite generating identity for all the displayed coefficients, so their calculation can be reproduced without enumerating 4^{16} full cubes. For $v \in \mathbb{Z}^3$, let

$$F_v(b) = [z^0 y^v] \prod_{\epsilon \in \{0,1\}^3} (zy^\epsilon + z^{-1}y^{-\epsilon} + bz^3y^{3\epsilon} + bz^{-3}y^{-3\epsilon}). \quad (\text{III.70.19})$$

The expressions are Laurent polynomials. Splitting the four-cube into its two three-faces gives the exact identity

$$P(b) = \sum_{v \in \mathbb{Z}^3} F_v(b) F_{-v}(b) = \sum_{v \in \mathbb{Z}^3} F_v(b)^2. \quad (\text{III.70.20})$$

Only $v \in [-12, 12]^3$ can occur, and $F_{-v} = F_v$ by negating all frequency choices. Multiplication of the eight four-term factors in (III.70.19), followed by (III.70.20), gives the table. This is an exact finite coefficient identity; the supplied verifier uses integer counts to reproduce every entry.

The first two coefficients also have the independent counts already proved above: $c_0 = 222$, and $c_1 = 16 \cdot 2 \cdot 27 = 864$, because there are sixteen positions for the unique third harmonic, two choices of its sign, and twenty-seven completions. Formula (III.70.16) gives $R_4(b) = P(b)/(16(1 + b^4)^4)$, proving (III.70.18). $\square$

Corollary III.70.8 (Explicit improved witness). *For any character of order greater than 24, put*

$$F = 7(\chi + \chi^{-1}) - (\chi^3 + \chi^{-3}).$$

Then

$$\|F\|_{U^2}^4 = 4804, \quad \|F\|_{U^4}^{16} = 5465198316262956, \quad (\text{III.70.21})$$

so

$$\frac{\|F\|_{U^4}^{16}}{\|F\|_{U^2}^{16}} = \frac{1366299579065739}{133153321267264} = 10.2611002569 \dots < \frac{111}{8}. \quad (\text{III.70.22})$$

The rescaled witness $F/16$ is bounded by one in absolute value. Consequently (III.70.22) proves the lower bound in Theorem III.70.1; one may take $G = \mathbb{Z}/29\mathbb{Z}$.

Proof. Substitute $b = -1/7$ in Proposition III.70.7 and multiply by the appropriate powers of seven. The bound $|F| \leq 16$ is immediate from the triangle inequality. Both normalized ratios are unchanged by rescaling. $\square$

[write] **6 October 2026; sources 28–33.** *A better witness, from the intake* (a review result, credited to the batch-121 intake, not to any source; recomputed at the write in exact integer arithmetic). On $\mathbb{Z}/7\mathbb{Z}$ the centered integer function $f = (181, 1000, -780, -932, 989, 540, -998)$ has

$$\frac{\|f\|_{U^4}^{16}}{\|f\|_{U^2}^{16}} = \frac{73456939652364023667491495159266974387074638671313075}{7160989329808722083386562416397488648766840301198728} = 10.2579317 \dots,$$

below the ratio $10.2611003 \dots$ of the witness here, so $c_4^{\text{odd}} \geq 0.5587725$ (exactly: the ratio is at most 0.5587725^{-4}). The group has order $7 < 24$, so congruence aliases occur; that is allowed, because c_4^{odd} is a supremum over all groups of odd order. Hence the two-harmonic witness is not optimal. The intake's

numerical optimization over centered functions on $\mathbb{Z}/N\mathbb{Z}$, $N = 5, \dots, 13$, found ratios down to about 10.25714 ($c_4^{\text{odd}} \gtrsim 0.558783$); that is evidence only.

[write] **6 October 2026; sources 34–39.** Source 34’s Proposition III.71.7, written the same day, reproduces this example (the same polynomial and witness) without citing this source; it is credited here and there.

For comparison, on $\mathbb{Z}/7\mathbb{Z}$ the same expression has $\|F\|_{U^2}^4 = 4804$ but $\|F\|_{U^4}^{16} = 5465448865508028$. Its ratio is $1366362216377007/133153321267264$, approximately 10.2615706719. This difference is a useful check on the characteristic hypothesis: the generic polynomial must not be substituted at a small character order. The verifier reproduces this second value by a modular coefficient expansion and, independently, by direct cube summation in $\mathbb{Q}[t]/(t^3 + t^2 - 2t - 1)$ with $t = \zeta_7 + \zeta_7^{-1}$.

III.70.5 The obstruction persists in every higher dimension

Theorem III.70.9 (Single characters are not local extremizers for $d \geq 4$). *Let $d \geq 4$ and let χ have order greater than $3 \cdot 2^{d-1}$. For*

$$f_{b,\theta} = e^{i\theta}\chi + e^{-i\theta}\chi^{-1} + b(e^{3i\theta}\chi^3 + e^{-3i\theta}\chi^{-3}),$$

the derivative at $b = 0$ of

$$R_d(b) = \frac{\|f_{b,\theta}\|_{U^d}^{2^d}}{\|f_{b,\theta}\|_{U^2}^{2^d}} \quad (\text{III.70.23})$$

is strictly positive. Hence every sufficiently small negative b strictly increases $\|f_{b,\theta}\|_{U^2}^4 / \|f_{b,\theta}\|_{U^d}^4$. In particular, a nonzero single real character pair of such order is not a local maximizer for the higher-order extremal problem.

Proof. Let $C_{d,j}$ count assignments $s : \{0, 1\}^d \rightarrow \{-3, -1, 1, 3\}$ having exactly j third harmonics and satisfying the integer equations

$$\sum_{\omega} s_{\omega} = 0, \quad \sum_{\omega} \omega_i s_{\omega} = 0 \quad (1 \leq i \leq d). \quad (\text{III.70.24})$$

The same half-cube argument as above shows that the character-order hypothesis turns all orthogonality congruences into these integer equations. The phase factor is then $e^{i\theta} \sum_{\omega} s_{\omega} = 1$. Thus

$$\|f_{b,\theta}\|_{U^d}^{2^d} = \sum_j C_{d,j} b^j, \quad R_d(b) = \frac{\sum_j C_{d,j} b^j}{[2(1 + b^4)]^{2^{d-2}}}.$$

All $C_{d,j}$ are nonnegative integers, and $C_{4,1} = 864$. There are pure-sign solutions in every dimension at least two, for example $s_{\omega} = (-1)^{\omega_1 + \dots + \omega_d}$, so $C_{d,0} > 0$. Place a solution with one third harmonic on one d -face and an independent pure-sign solution on the other. Their separate moment conditions imply all the $(d+1)$ -dimensional conditions. Either face can contain the unique third harmonic, and the resulting assignments are distinct. Therefore

$$C_{d+1,1} \geq 2C_{d,1}C_{d,0} > 0. \quad (\text{III.70.25})$$

It follows that $R'_d(0) = C_{d,1}/2^{2^{d-2}} > 0$. The desired inverse ratio is $R_d(b)^{-1/2^{d-2}}$, which decreases strictly with $R_d(b)$, proving the assertion. $\square$

[write] **6 October 2026; sources 28–33.** *Answers two questions.* Together with Proposition III.70.2 and the witness above, this theorem answers, for $d \geq 4$, the sub-question of source 03’s Research question V.14.1 whether the odd-order constant is attained by real parts of characters: no (for $d = 3$ it is, by source 04’s Theorem III.3.4). It also answers the second sentence of source 29’s Question V.117.1 (“Is a single Fourier inverse pair extremal?”): no. Item III.1.

[write] **6 October 2026; sources 34–39.** At $d = 4$ source 34 gives a global form: every function supported on one inverse pair of a character of odd order at least five has ratio at most $(8/111)^{1/4}$, below the witness above (Theorem III.71.1 and the note there).

[write] **7 October 2026; sources 40–46.** Source 45 also re-derives the single-pair non-extremality at $d = 4$ (its Section III.90, with the dual identity and order three), without citing this source or source 34.

III.70.6 What is settled and what remains open

The proved bound (III.70.3) improves the prior general comparison for every $d \geq 4$. The exact examples give a stronger lower bound for the fourth-order constant, and Theorem III.70.9 rules out extending the third-order single-character equality model to higher order in general. These conclusions do not identify the sharp value of c_4^{odd} , establish that two harmonics suffice at an optimum, or classify maximizing sequences. The recurrence (III.70.25) proves the sign of a first variation; it does not determine an optimal higher-dimensional constant.

A natural next problem is to optimize the two complementary bounds in Lemma III.70.5, retaining the separate fourth-power masses at χ^3 and χ^5 . A more substantial improvement would use dual functions of a two-harmonic trial function instead of a cosine, together with compatible concentration estimates. Another precise question is whether the global fourth-order supremum is attained on a finite cyclic group or requires a limiting spectral profile. Any such investigation must keep small-order aliases separate from integer frequency relations.

The exact quartic-count consequence (III.70.14) is local. It does not by itself improve a final Szemerédi bound. The supplied standard-library program verifies the finite generating identities, the two explicit finite-group examples, the rational comparisons in the universal estimate, and a finite set of the constructive higher-dimensional certificates. The general proofs are those above; no numerical optimization or proof-assistant verification is claimed as a substitute.

III.71 A stronger fourth-dimensional norm comparison

Source 34, delivered Section 10.

The torsion analysis and the analytic norm comparison address different losses in a centered cube expansion. We now improve the analytic coefficient. Throughout this section, G is a finite abelian group of odd order, physical averages are normalized, and Fourier sums use counting measure. In this section we also use additive notation for frequency labels: $-\xi$ denotes ξ^{-1} and $j\chi$ denotes χ^j . For real, centered, nonzero f , recall

$$c_d^{\text{odd}} = \sup_{G,f} \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4}. \quad (\text{III.71.1})$$

The merged repository report proves $c_3^{\text{odd}} = 2^{-1/2}$, with equality precisely for one conjugate Fourier pair [155]. The subsequently received moment manuscript already improves the fourth-dimensional




Figure III.70.1: The exact polynomial ratio $R_4(b)$ in the two-harmonic family. The marked point is the rational certificate $b = -1/7$. The curve illustrates the proved local improvement over $b = 0$; it is not a computation of the global extremal constant.

bracket to

$$(8/111)^{1/4} \leq c_4^{\text{odd}} \leq (8/35)^{1/4}, \quad (\text{III.71.2})$$

whose endpoints are approximately 0.518133595 and 0.691441569 [156]. That later manuscript is the relevant baseline for the following improvement.

Theorem III.71.1 (Improved fourth-dimensional bracket). *Every real centered function on every finite abelian group of odd order satisfies*

$$\|f\|_{U^2}^4 \leq \frac{2}{3} \|f\|_{U^4}^4. \quad (\text{III.71.3})$$

More precisely,

$$\frac{4804}{5465198316262956^{1/4}} \leq c_4^{\text{odd}} \leq \frac{2}{3}. \quad (\text{III.71.4})$$

The lower endpoint is approximately 0.558729391. It is attained as a lower-bound certificate by a function with two conjugate Fourier pairs on $\mathbb{Z}/29\mathbb{Z}$. In contrast, every function supported on one conjugate Fourier pair has ratio at most $(8/111)^{1/4}$.

[write] **6 October 2026; sources 34–39.** Where this stands (checked at the write). In source 29’s normalization $\|f\|_{U^4}^{16} \geq C \|f\|_{U^2}^{16}$ the theorem is $C = (3/2)^4 = 81/16$, so $c_4^{\text{odd}} = \sup C^{-1/4}$ -bound $\leq 2/3$; source 29 has $C = 35/8$ (0.6914416), source 35, written the same day from the same pin and not cited here, $C = 5$ ($5^{-1/4} = 0.6687403$; Section III.72), and source 28 $(323/160)^2$. So $2/3$ is the best upper bound among sources 01–39. It is not the best known: source 45, placed in this directory by **2dfe46f44** and not yet written in, records $C = 4795/832$, that is $c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$ (its placement, checked at intake; for the write that adds it). The lower endpoint here is source 28’s witness (0.5587294, Corollary III.70.8), improved by the batch-121 intake’s exact certificate on $\mathbb{Z}/7\mathbb{Z}$ to 0.5587725 (note there). Among sources 01–39 the bracket is therefore $0.5587725 \leq c_4^{\text{odd}} \leq 2/3$ (item III.1). The statement that one-pair functions have ratio at most $(8/111)^{1/4} = 2/222^{1/4}$ is the global form, at $d = 4$ and every odd order at least five, of source 28’s non-extremality theorem; for order three see Lemma III.71.4.

[write] **7 October 2026; sources 40–46.** *Superseded by source 45* (written the same day, now printed in Section III.89): its Theorem III.90.1 has the coefficient $4795/832 = 5.7632212 > 81/16$, so $c_4^{\text{odd}} \leq 0.6454070$ and the bracket among sources 01–46 is $0.5587725 \leq c_4^{\text{odd}} \leq 0.6454070$. The statement “best among sources 01–39” above stays true for those sources. Source 45 also re-derives the single-pair results of this section without citing it (note in Section F.50).

[write] **7 October 2026; sources 47–53.** *Superseded again by source 49* (Section III.92): coefficient $29/4$, so $0.5587725 \leq c_4^{\text{odd}} \leq 0.6094183$ among sources 01–53. Source 49 builds its dual test on this section’s two-harmonic example and cites this source for the balanced-moment profile and the bound $2/3$.

The theorem improves both endpoints of the audited bracket. In particular, the one-pair extremizers in dimension three cease to be extremizers in dimension four. Neither endpoint in (III.71.4) is asserted to be the optimum.

III.71.1 Retaining the dominant-pair moment profile

Write $a_\xi = \widehat{f}(\xi)$, so $a_{-\xi} = \overline{a_\xi}$ and $a_0 = 0$. Set

$$S = \sum_{\xi} |a_\xi|^4 = \|f\|_{U^2}^4, \quad p_j = \frac{2|a_{\xi_j}|^4}{S}, \quad p = \max_j p_j,$$

where one ξ_j is chosen from each nonzero inverse pair. Thus $\sum_j p_j = 1$. The incoming moment argument [156] begins with the positive Fourier expansion of the autocorrelation

$$c_f(h) = \mathbb{E}_x f(x)f(x-h) = \sum_{\xi} |a_\xi|^2 \xi(h).$$

We reproduce that argument and retain the additional information needed to combine it with a dual-function estimate.

Lemma III.71.2 (A profile for a specified Fourier pair). *For $m \geq 1$, put*

$$C_m = 2^{-m} \binom{2m}{m}, \quad H_m(p) = \sum_{k=0}^m \binom{m}{k}^2 p^k (1-p)^{m-k}.$$

For any selected pair of mass p , one has

$$\mathbb{E}_h c_f(h)^{2m} \geq C_m S^m H_m(p). \quad (\text{III.71.5})$$

Consequently, when $m = 2^{d-2}$,

$$\frac{\|f\|_{U^d}^4}{S} \geq (C_m H_m(p))^{1/m}. \quad (\text{III.71.6})$$

For $m = 4$ this gives

$$\frac{\|f\|_{U^4}^4}{S} \geq A(p) := \left[\frac{35}{8} \{1 + 12p(1-p) + 6p^2(1-p)^2\} \right]^{1/4}. \quad (\text{III.71.7})$$

If the largest pair has $p \leq 1/2$, one additionally has

$$\frac{\|f\|_{U^4}^4}{S} \geq (1015/64)^{1/4}. \quad (\text{III.71.8})$$

Proof. Put $b_j = |a_{\xi_j}|^2$, so $S = 2 \sum_j b_j^2$. In the Fourier expansion of c_f^{2m} , every coefficient is nonnegative. Retain only monomials using each ξ_j and $-\xi_j$ equally often. Writing their common multiplicity as k_j gives

$$\mathbb{E} c_f^{2m} \geq C_m S^m \sum_{\sum_j k_j = m} \binom{m}{k_1, k_2, \dots}^2 \prod_j p_j^{k_j}. \quad (\text{III.71.9})$$

This is the positive balanced-monomial step of [156]. Group its terms according to the multiplicity k of the specified pair. The remaining multinomial coefficient is an integer at least one, so its square is at least itself. The multinomial theorem therefore bounds the contribution of the other pairs below by $(1-p)^{m-k}$. This proves (III.71.5).

The recursive norm identity and monotonicity give

$$\|f\|_{U^d}^{2^d} = \mathbb{E}_h \|f(\cdot) f(\cdot - h)\|_{U^{d-1}}^{2^{d-1}} \geq \mathbb{E}_h |c_f(h)|^{2^{d-1}}.$$

Since $2^{d-1} = 2m$ and $S^m = \|f\|_{U^2}^{2^d}$, this proves (III.71.6). Direct expansion gives $H_4(p) = 1 + 12p(1-p) + 6p^2(1-p)^2$.

For the last claim, interpret the normalized balanced sum in (III.71.9) as the expected multinomial coefficient under the multinomial distribution of four trials with probabilities p_j . That coefficient is one for a pure composition and at least four for every mixed composition. Hence the sum is at least $4 - 3 \sum_j p_j^4$. Because $\sum_j p_j^4 \leq p^3 \sum_j p_j \leq 1/8$, the lower bound is $29/8$. Multiplication by $C_4 = 35/8$ proves (III.71.8). $\square$

Thus near the old moment endpoint, almost all Fourier fourth-power mass lies in one pair. The next step uses information that the moment estimate alone discarded: a one-pair function has a substantially larger fourth Gowers norm than that endpoint permits.

III.71.2 The exact fourth dual function of a real character

For real h , define

$$\mathcal{D}_4 h(x) = \mathbb{E}_{t_1, t_2, t_3, t_4} \prod_{\substack{v \in \{0,1\}^4 \\ v \neq 0}} h(x + v \cdot t).$$

The Gowers–Cauchy–Schwarz inequality supplies the dual certificate

$$|\mathbb{E} f \mathcal{D}_4 h| \leq \|f\|_{U^4} \|h\|_{U^4}^{15}, \quad \mathbb{E} h \mathcal{D}_4 h = \|h\|_{U^4}^{16}. \quad (\text{III.71.10})$$

The general use of a dual function is standard. The useful input here is its exact value on a real character, including the two additional harmonics.

Lemma III.71.3 (The fifteen-vertex dual polynomial). *Let χ have odd order at least five, let $\theta \in \mathbb{R}$, and put*

$$\phi_j = e^{ij\theta} \chi^j + e^{-ij\theta} \chi^{-j}.$$

Then

$$\mathcal{D}_4 \phi_1 = 111\phi_1 + 27\phi_3 + \phi_5, \quad \|\phi_1\|_{U^4}^{16} = 222. \quad (\text{III.71.11})$$

These identities hold for arbitrary θ , including when the character order is five or seven.

Proof. Expand the fifteen factors into their two character terms. For each coordinate, orthogonality requires a sum of eight signs to vanish modulo the order of χ . This sum is even and has absolute value at most eight. Since that order is odd and at least five, the sum must vanish as an integer.

The positive-sign subset T of the fifteen nonorigin vertices therefore has exactly four ones in every coordinate. If $|T| = k$, its output harmonic is $2k - 15$. The complete count is

$$\begin{array}{c|cccccc} k & 5 & 6 & 7 & 8 & 9 & 10 \\ \hline \#T & 1 & 27 & 111 & 111 & 27 & 1. \end{array} \quad (\text{III.71.12})$$

Here is an elementary verification. Replace each selected vertex by its set of zero coordinates, a proper subset of $\{1, 2, 3, 4\}$. For five selected vertices each coordinate must occur once among these zero sets. The only five distinct sets with total size four are the empty set and the four singletons, giving one choice.

For six selected vertices every coordinate has degree two. If the empty set is absent, the four singleton sets and a perfect matching are forced, giving three choices. If it is present, either the three pair sets form a Hamilton path, with singleton sets at the two endpoints (twelve choices), or there are one triple, one pair, and three singletons. In the latter case the vertex without a singleton belongs to both the pair and triple, and the other endpoint of the pair is the omitted vertex of the triple. There are $4 \cdot 3 = 12$ choices. The total is 27.

For seven selected vertices every coordinate has degree three. Let $z \in \{0, 1\}$ record the empty set, and let u, e, t be the numbers of singleton, pair, and triple sets. The complete list is

z	t	u	e	number
0	0	2	5	6
0	1	3	3	$4 + 24$
0	2	4	1	6
1	0	0	6	1
1	1	1	4	$12 + 12$
1	2	2	2	$18 + 24$
1	3	3	0	4

To verify completeness, the degree and cardinality equations give $e = 5 + z - 2t$ and $u = 2 - 2z + t$. If T_0 is the set of vertices omitted by the triples and U_0 is the set of singleton vertices, the graph formed by the pairs must have degree $3 - t + \mathbf{1}_{T_0}(v) - \mathbf{1}_{U_0}(v)$ at v . The rows count, respectively: a missing edge in K_4 ; a star or a Hamilton path with prescribed internal vertices; a prescribed edge; K_4 ; a four-cycle or the complement of a two-edge path; a perfect matching or a two-edge path; and equal sets $T_0 = U_0$. The entries sum to 111.

Complementation within all fifteen vertices preserves the four coordinate counts and exchanges k with $15 - k$. Sizes at most four are impossible: four distinct vertices cannot all have a one in every coordinate. This proves (III.71.12), hence the dual identity. Finally, character orthogonality in $\mathbb{E}\phi_1\mathcal{D}_4\phi_1$ leaves just 222. In orders five and seven the third and fifth harmonics still do not overlap the first; in order five ϕ_5 is constant, which causes no change. $\square$

[write] **6 October 2026; sources 34–39.** The identity and the value 222 are also source 28’s Lemma III.70.3, proved there by a face-splitting count; source 35’s Lemma III.72.3 recounts them by a third decomposition. All three sources were written on the same day; none cites another for this identity.

The order-three case requires its own calculation, since additional sign sums can vanish modulo three.

Lemma III.71.4 (The order-three endpoint). *For a character χ of order three and $\phi = e^{i\theta}\chi + e^{-i\theta}\chi^{-1}$, let $r = 2\cos(3\theta)$. Then*

$$\|\phi\|_{U^4}^{16} = 270 - 16r^2 + 8r^4 = 262 + 8(r^2 - 1)^2 \geq 262. \quad (\text{III.71.13})$$

If $\widehat{f}(\chi) = \rho e^{i\theta}$, then $\|f\|_{U^4} \geq \rho\|\phi\|_{U^4}$.

Proof. The three values a, b, c of ϕ satisfy $a + b + c = 0$, $ab + bc + ca = -3$, and $abc = r$. Flipping cube coordinates shows that the average associated with a step vector depends only on its number of nonzero coordinates. Their multiplicities are 1, 8, 24, 32, 16. Consequently, after also averaging the basepoint,

$$243\|\phi\|_{U^4}^{16} = \sum a^{16} + 8 \sum (ab)^8 + 24r^4 \sum a^4 + 32r^4 \sum (ab)^2 + 16r^5 \sum a,$$

where each sum has three cyclic terms. Newton recurrences give

$$\sum a^4 = 18, \quad \sum (ab)^2 = 9, \quad \sum a^{16} = 13122 + 11664r^2 + 360r^4,$$

and $\sum (ab)^8 = 6561 - 1944r^2 + 108r^4$. Substitution proves (III.71.13).

Conditional expectation onto the quotient defined by χ is averaging over translations in its kernel, hence is a contraction for the U^4 norm. Since f is centered, this conditional expectation is exactly $\rho\phi$. This proves the last claim. $\square$

III.71.3 Coupling the moment and dual certificates

Proposition III.71.5 (A second lower bound from the largest pair). *Put*

$$K = \frac{222^{1/4}}{2}, \quad b = \frac{82^{3/4}}{111}, \quad B(p) = K[p^{1/4} - b(1-p)^{1/4}]_+^4,$$

where $[x]_+ = \max(x, 0)$. On every odd-order group,

$$\frac{\|f\|_{U^4}^4}{S} \geq B(p). \quad (\text{III.71.14})$$

Proof. Choose a largest pair and write $a_\chi = \rho e^{i\theta}$, where $\rho = (pS/2)^{1/4}$. If its order is at least five, Lemma III.71.3 and centering give

$$\mathbb{E}f\mathcal{D}_4\phi_1 \geq 222\rho - 54|a_{3\chi}| - 2|a_{5\chi}|.$$

The nonzero pairs $\{\pm 3\chi\}$ and $\{\pm 5\chi\}$ are distinct from each other and from $\{\pm\chi\}$. If $5\chi = 0$, its coefficient vanishes. Hölder therefore yields

$$\begin{aligned} 54|a_{3\chi}| + 2|a_{5\chi}| &\leq (54^{4/3} + 2^{4/3})^{3/4}(|a_{3\chi}|^4 + |a_{5\chi}|^4)^{1/4} \\ &\leq 2 \cdot 82^{3/4} (S(1-p)/2)^{1/4}. \end{aligned}$$

Apply (III.71.10), using $\|\phi_1\|_{U^4}^{16} = 222$, and raise to the fourth power. This gives precisely (III.71.14).

If the selected character has order three, Lemma III.71.4 instead gives

$$\frac{\|f\|_{U^4}^4}{S} \geq \frac{262^{1/4}}{2} p \geq Kp \geq B(p).$$

Thus no group-exponent restriction has been omitted. $\square$

The two separate profiles already improve the incoming bound, but their interaction gives a larger gain. The third and fifth harmonics that weaken the dual certificate also force additional *positive* monomials in the autocorrelation moment.

Lemma III.71.6 (The resonance cost of the dual error). *Suppose the selected character χ has odd order at least five. Define*

$$q_3 = \frac{2|a_{3\chi}|^4}{S}, \quad q_5 = \frac{2|a_{5\chi}|^4}{S}, \quad L = 27q_3^{1/4} + q_5^{1/4},$$

with $q_5 = 0$ when $5\chi = 0$. Then, writing $R = \|f\|_{U^4}^4/S$, one has simultaneously

$$R \geq K[p^{1/4} - L/111]_+^4, \tag{III.71.15}$$

$$R^4 \geq \frac{35}{8}H_4(p) + p^{7/2}(21\sqrt{q_3} + 7\sqrt{q_5}) \geq \frac{35}{8}H_4(p) + \frac{7}{244}p^{7/2}L^2. \tag{III.71.16}$$

Proof. The first inequality is the proof of Proposition III.71.5 before applying Hölder to the two unwanted harmonics. For the second, return to the expansion of c_f^8 . In addition to all pair-balanced monomials retained in Lemma III.71.2, retain those containing one 3χ , two χ , and five $-\chi$, and their inverses. Their total contribution is

$$2\frac{8!}{5!2!}|a_\chi|^{14}|a_{3\chi}|^2 = 336|a_\chi|^{14}|a_{3\chi}|^2.$$

Likewise, the monomials containing one 5χ , one χ , and six $-\chi$, together with their inverses, contribute $112|a_\chi|^{14}|a_{5\chi}|^2$. These are valid zero-sum frequency tuples by the integer relations $3 + 2 - 5 = 0$ and $5 + 1 - 6 = 0$. They are disjoint from the pair-balanced monomials and from each other: each uses exactly one occurrence from a different nonzero pair. In order five the second family has zero coefficient; in order seven the pairs remain distinct. Any further modular relations only add nonnegative terms.

Divide by S^4 and use $|a_\chi|^{14}|a_{j\chi}|^2 = (S/2)^4 p^{7/2} \sqrt{q_j}$. The moment lower bound for R^4 follows. Finally, weighted Cauchy–Schwarz gives

$$(27q_3^{1/4} + q_5^{1/4})^2 \leq \left(\frac{27^2}{21} + \frac{1}{7}\right)(21\sqrt{q_3} + 7\sqrt{q_5}) = \frac{244}{7}(21\sqrt{q_3} + 7\sqrt{q_5}),$$

which proves the last inequality. $\square$

Proof of the upper bound in Theorem III.71.1. The zero function is immediate, so suppose $S > 0$ and put $R = \|f\|_{U^4}^4/S$. If $p \leq 1/2$, (III.71.8) gives $R > 3/2$. On $[1/2, 1]$, A is decreasing, and for $p_0 = 49/50$ exact arithmetic gives

$$A(p_0)^4 - \frac{81}{16} = \frac{1757921}{5000000} > 0. \tag{III.71.17}$$

Hence $1/2 \leq p \leq p_0$ also implies $R > 3/2$.

It remains to consider $p \geq p_0$. If the selected character has order three, Lemma III.71.4 gives $R \geq Kp > (193/100)(49/50) > 3/2$. Suppose instead that its order is at least five and, towards a contradiction, $R \leq 3/2$. The following rational comparisons are elementary:

$$p^{1/4} > \frac{99}{100}, \quad K > \frac{193}{100}, \quad \left(\frac{3}{2K}\right)^{1/4} < \frac{47}{50}.$$

Indeed $(99/100)^4 < 49/50$, $(193/100)^4 < 222/16$, and $(47/50)^4 > 150/193$. By (III.71.15), these comparisons force

$$\frac{L}{111} > \frac{99}{100} - \frac{47}{50} = \frac{1}{20}, \quad L > \frac{111}{20}.$$

Since $H_4(p) \geq 1$ and $p^{7/2} \geq p^4 \geq (49/50)^4$, (III.71.16) now yields

$$\begin{aligned} R^4 - \frac{81}{16} &> \frac{35}{8} + \frac{7}{244} \left(\frac{49}{50}\right)^4 \left(\frac{111}{20}\right)^2 - \frac{81}{16} \\ &= \frac{77821791847}{610000000000} > 0, \end{aligned} \tag{III.71.18}$$

contradicting $R \leq 3/2$. Thus $R > 3/2$ for nonzero f , which proves the stated uniform inequality. $\square$

III.71.4 An explicit two-pair improvement and the variational obstruction

Proposition III.71.7 (An exact two-harmonic polynomial). *On $\mathbb{Z}/29\mathbb{Z}$, let $\chi(x) = e^{2\pi ix/29}$ and*

$$f_t = \chi + \chi^{-1} + t(\chi^3 + \chi^{-3}), \quad t \in \mathbb{R}.$$

Then $\|f_t\|_{U^2}^4 = 2(1 + t^4)$ and

$$\begin{aligned} P(t) := \|f_t\|_{U^4}^{16} &= 222 + 864t + 4080t^2 + 9120t^3 + 25640t^4 \\ &\quad + 31808t^5 + 55120t^6 + 49280t^7 + 52360t^8 \\ &\quad + 28320t^9 + 24160t^{10} + 4256t^{11} + 5920t^{12} + 222t^{16}. \end{aligned} \tag{III.71.19}$$

In particular $7^{16}P(-1/7) = 5465198316262956$, which proves the lower bound in (III.71.4).

Proof. The U^2 identity is Fourier orthogonality. For the second identity, split the four-cube into its two three-dimensional faces. An expanded term assigns a frequency in $\{-3, -1, 1, 3\}$ to each vertex. Orthogonality in the fourth step and the basepoint requires the total frequency on each face to be zero. These and all remaining coordinate conditions hold over the integers: each relevant sum has eight terms and absolute value at most $24 < 29$.

For $v \in \{0, 1\}^3$ assign a frequency r_v , and define

$$A_m(t) = \sum_{\substack{r_v \in \{-3, -1, 1, 3\} \\ \sum_v r_v = 0, \sum_v r_v v = m}} t^{\#\{v: |r_v|=3\}}, \quad m \in \mathbb{Z}^3.$$

The exact finite identity is

$$P(t) = \sum_{m \in \mathbb{Z}^3} A_m(t) A_{-m}(t). \tag{III.71.20}$$

It supplies a short exhaustive integer certificate for (III.71.19): enumerate the 4^8 assignments on one face, discard those with nonzero total, and group the others by their three integer moments and their count of frequencies of absolute value three. Multiplying opposite-moment polynomials gives the coefficients displayed in (III.71.19). The accompanying standard-library verification program implements exactly this finite sum, using integers and rational arithmetic only. Thus the polynomial certificate requires neither sampling nor numerical optimization.

Finally,

$$\frac{\|f_{-1/7}\|_{U^2}^4}{\|f_{-1/7}\|_{U^4}^4} = \frac{2(1 + 7^{-4})}{P(-1/7)^{1/4}} = \frac{4804}{5465198316262956^{1/4}}.$$

$\square$

[write] **6 October 2026; sources 34–39.** *This is source 28’s example* (Corollary III.70.8): the same functions on $\mathbb{Z}/29\mathbb{Z}$, the same polynomial and the same witness $7^{16}P(-1/7) = 5465198316262956$. The intake found that the best member of the family, at $t \approx -0.145611$, gives 0.5587462, and the batch-121 intake’s certificate on $\mathbb{Z}/7\mathbb{Z}$ gives 0.5587725 (note after source 28’s corollary); neither is claimed by a source as optimal.

The full polynomial is useful for an explicit certificate, but the failure of one-pair extremality follows already from the dual function. At $t = 0$, its third harmonic gives $P'(0) = 16 \cdot 2 \cdot 27 = 864 > 0$, whereas $\frac{d}{dt}\|f_t\|_{U^2}^4|_{t=0} = 0$. Therefore small negative t strictly increases the ratio. This mechanism persists for arbitrary initial phase and every character order at least five: the third harmonic lies outside the original pair and is nonzero.

To compare with every one-pair model, if $f = \rho(e^{i\theta}\chi + e^{-i\theta}\chi^{-1})$ then $\|f\|_{U^2}^4 = 2\rho^4$. For character order at least five, Lemma III.71.3 gives ratio $2/222^{1/4}$; for order three, Lemma III.71.4 gives ratio at most $2/262^{1/4}$, which is smaller. Thus the two-pair certificate exceeds every one-pair model on every finite odd-order group.

Remark III.71.8 (What a genuine optimizer must satisfy). On a fixed finite group, minimize $\|f\|_{U^4}^{16}$ subject to $\|f\|_{U^2}^4 = 1$ and $\mathbb{E}f = 0$. A minimizer exists by finite-dimensional compactness. The Lagrange-multiplier equation is

$$\mathcal{D}_4 f = \lambda \mathcal{D}_2 f + \mu, \quad \widehat{\mathcal{D}_2 f}(\xi) = |a_\xi|^2 a_\xi,$$

for real constants λ, μ . The extra third harmonic in (III.71.11) prevents a one-pair function of character order at least five from satisfying this equation. Order three is an exception to that stationarity statement, since its centered function space already consists of one Fourier pair.

III.71.5 The resulting fourth-dimensional cube coefficient

Proposition III.81.2, recalling the established odd-order centered expansion [155, 156], gives

$$Q_4(\delta + f) = \delta^{16} + 100\delta^{12}\|f\|_{U^2}^4 + R, \quad |R| \leq \sum_{j=5}^{16} \binom{16}{j} \delta^{16-j} \|f\|_{U^4}^j,$$

for $\delta \geq 0$. The theorem therefore gives the following direct refinement of the coefficient used at Lemma 3.10.

Corollary III.71.9 (A smaller leading coefficient). *Let $u = \|f\|_{U^4}$. Under the preceding hypotheses,*

$$Q_4(\delta + f) \leq \delta^{16} + \frac{200}{3}\delta^{12}u^4 + \sum_{j=5}^{16} \binom{16}{j} \delta^{16-j} u^j. \quad (\text{III.71.21})$$

The coefficient $200/3 \approx 66.666667$ replaces the audited incoming coefficient $100(8/35)^{1/4} \approx 69.144157$.

[write] **6 October 2026; sources 34–39.** The leading coefficient at $d = 4$ is now $200/3 = 66.667$ (against $100 \cdot 5^{-1/4} = 66.874$ for source 35, $100(8/35)^{1/4} = 69.144$ for source 29, 70.382 for source 28 and $100/\sqrt{2} = 70.711$ from source 04’s comparison); recomputed at the write.

If $u^{16} \leq \alpha$, the leading correction in (III.71.21) is $(200/3)\delta^{12}\alpha^{1/4}$ and the remaining terms have powers $\alpha^{j/16}$, $j \geq 5$. The two-pair example implies that the optimal universal leading coefficient is at least $100 \cdot 4804/5465198316262956^{1/4} \approx 55.872939$: scale that fixed example to zero amplitude while keeping a positive density fixed. At sufficiently small amplitude it also gives a weight in $[0, 1]$. This is a local counting improvement; the subsequent density-increment losses have not been recomputed.

Remark III.71.10 (The same coupling in a few higher dimensions). For $m = 2^{d-2}$, Lemma III.71.2 and monotonicity $\|f\|_{U^d} \geq \|f\|_{U^4}$ combine the profiles $(C_m H_m(p))^{1/m}$ and $B(p)$. The identity

$$H_m(p) = \sum_{j=0}^{\lfloor m/2 \rfloor} \binom{m}{2j} \binom{2j}{j} \{p(1-p)\}^j$$

follows by extracting the coefficient of z^m in $(\sqrt{p(1-p)} + z + \sqrt{p(1-p)}z^2)^m$; it shows that the first profile decreases on $[1/2, 1]$. When $p \leq 1/2$, the balanced-composition argument gives the stronger bound

$$\frac{\|f\|_{U^d}^4}{S} \geq [C_m \{m - (m-1)2^{-(m-1)}\}]^{1/m}.$$

For $m = 4, 8, 16, 32, 64$, this last bound exceeds K , while $C_m^{1/m} < K$. Both comparisons become rational inequalities after raising to the m th power; the accompanying certificate checks them exactly. For each $d \in \{5, 6, 7, 8\}$, let p_d be the unique point of $[1/2, 1]$ satisfying

$$(C_m H_m(p_d))^{1/m} = B(p_d), \quad \beta_d = (C_m H_m(p_d))^{-1/m}.$$

The decreasing and increasing profiles meet once: their endpoint order is opposite, and the diffuse branch is stronger than K . Their maximum is at least their common value at p_d , proving $c_d^{\text{odd}} \leq \beta_d$. The following are rounded approximations to these exact implicit bounds:

d	incoming moment bound	coupled bound
5	0.612819643	0.609912380
6	0.565388162	0.564904412
7	0.537413970	0.537382843
8	0.5211683364	0.5211682954

These are consequences of the same two proved profiles, not applications of the weaker uniform bound $2/3$. The fourth dual certificate alone does not improve the incoming moment bound in dimensions $d \geq 9$.

[write] **6 October 2026; sources 34–39.** Recomputed at intake by bisection: the four coupled bounds and the four bounds of source 29 are reproduced to the printed digits; the gain over source 29's B_d is about $3 \cdot 10^{-3}$ at $d = 5$ and $4 \cdot 10^{-8}$ at $d = 8$, and none for $d \geq 9$, as stated. So for $5 \leq d \leq 8$ the best upper bounds among sources 01–39 are the β_d here; for $d \geq 9$ they are source 29's B_d (Corollary III.69.1).

[write] **7 October 2026; sources 47–53.** Source 49's Theorem III.93.1 gives $c_5^{\text{odd}} \leq 29/50 = 0.58$, below $\beta_5 = 0.6099124$ here; for $d = 6, 7, 8$ the coupled bounds β_d of this remark remain the best (they are below 0.58).

III.72 A stronger gap between real centered uniformity norms

Source 35, delivered Section 6.

Delivered part heading of source 35: A stronger real fourth-order uniformity gap

We use normalized counting measure on a finite abelian group G : $\mathbb{E}_x u(x) = |G|^{-1} \sum_{x \in G} u(x)$. Let $\mathcal{C}$ denote complex conjugation, so $\mathcal{C}z = \bar{z}$. For $f : G \rightarrow \mathbb{C}$, the normalized Gowers quantity is

$$\|f\|_{U^d(G)}^{2^d} = \mathbb{E}_{x, h_1, \dots, h_d \in G} \prod_{\omega \in \{0,1\}^d} \mathcal{C}^{|\omega|} f(x + \omega \cdot h), \quad \omega \cdot h = \sum_{i=1}^d \omega_i h_i. \quad (\text{III.72.1})$$

This defines a norm for $d \geq 2$; at $d = 1$ it is the seminorm $|\mathbb{E}f|$. For a character $\xi \in \widehat{G}$, our Fourier normalization is

$$\widehat{f}(\xi) = \mathbb{E}_{x \in G} f(x) \overline{\xi(x)}. \quad (\text{III.72.2})$$

We suppress G in the norm subscript when the ambient group is fixed. For real functions the conjugations in (III.72.1) have no effect.

The original Gowers inequalities compare successive uniformity norms by monotonicity [1]. Realness and centering permit a strict comparison on odd-order groups. The earlier research compilation proves

$$\|f\|_{U^3}^8 \geq 2 \|f\|_{U^2}^8$$

for real mean-zero functions, with equality precisely for a single conjugate Fourier pair [159]. The subsequent manuscript *Logarithmic Selection, Local Freiman Models, and Higher Uniformity* shows that this equality model does not persist at order four, and obtains

$$\|f\|_{U^d}^4 \geq \sqrt{\frac{323}{160}} \|f\|_{U^2}^4 \quad (d \geq 4)$$

[162]. More recently, the positive-autocorrelation moment argument of [163] gives the stronger fourth-order bound

$$\|f\|_{U^4}^{16} \geq \frac{35}{8} \|f\|_{U^2}^{16}.$$

We increase its coefficient from $35/8$ to 5 . The additional information is a positive term counting three-term progressions in the Fourier amplitude distribution. Its third-harmonic specialization controls the principal negative direction in the fourth-order dual function of a cosine.

Theorem III.72.1 (A stronger real centered fourth-order norm gap). *Let G be a finite abelian group of odd order, and let $f : G \rightarrow \mathbb{R}$ have mean zero. Then*

$$\|f\|_{U^4}^{16} \geq 5 \|f\|_{U^2}^{16}. \quad (\text{III.72.3})$$

In fact (III.72.3) is strict for nonzero f . Consequently, for every integer $d \geq 4$,

$$\|f\|_{U^d}^4 \geq 5^{1/4} \|f\|_{U^2}^4. \quad (\text{III.72.4})$$

The constant 5 is not asserted to be optimal. The improvement over the strongest inspected prior estimates is at $d = 4$.

[write] 6 October 2026; sources 34–39. *Superseded the same day by source 34.* Source 34’s Theorem III.71.1, written from the same pin and not cited here, gives the coefficient $81/16 = 5.0625 > 5$, that is $c_4^{\text{odd}} \leq 2/3$ against $5^{-1/4} = 0.6687403$ here; source 45 (placed, not yet written) records $4795/832$. The statement that 5 improves “the strongest inspected U^4 coefficient $35/8$ ” was true at the pin. The energy inequality of Theorem III.72.2 is a separate result that source 34 does not contain. The intake re-derived the proof and recomputed every rational comparison of the scalar certificate exactly.

Define the extremal constants

$$c_d^{\text{odd}} = \sup_{\substack{G \text{ finite abelian, } |G| \text{ odd} \\ f: G \rightarrow \mathbb{R}, \mathbb{E}f=0, f \neq 0}} \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4}. \quad (\text{III.72.5})$$

The new fourth-order bound gives

$$c_4^{\text{odd}} \leq 5^{-1/4} = 0.668740304976 \dots, \quad (\text{III.72.6})$$

improving $(8/35)^{1/4} = 0.691441569\dots$ from [163], and also the earlier $\sqrt{160/323} = 0.703815343078\dots$. The known rational lower example from [162] consequently gives the smaller interval

$$\left(\frac{133153321267264}{1366299579065739}\right)^{1/4} \leq c_4^{\text{odd}} \leq 5^{-1/4}. \quad (\text{III.72.7})$$

The left endpoint is approximately 0.558729390985. We import that example only for comparison; the proof of Theorem III.72.1 does not use it.

The propagation to $d \geq 5$ in (III.72.4) is valid, but is weaker than the prior central-binomial estimate at those orders. The combined statement, with $q = 2^{d-2}$, is

$$c_d^{\text{odd}} \leq \min \left\{ 5^{-1/4}, \frac{2}{\binom{2q}{q}^{1/q}} \right\} \quad (d \geq 4). \quad (\text{III.72.8})$$

The second entry is the bound of [163]; at $d = 5$ it is approximately 0.612819643. Thus no new higher-dimensional record is inferred merely from monotonicity.

III.72.1 The Fourier identity and the retained energy terms

All averages use normalized counting measure. We write the dual group additively and put

$$a_r = \widehat{f}(r), \quad \lambda_r = |a_r|^2, \quad S = \sum_r \lambda_r^2 = \|f\|_{U^2}^4, \quad M = \sum_r \lambda_r^4.$$

Thus $a_{-r} = \overline{a_r}$ and $a_0 = 0$. The sums over the dual group are ordinary, unnormalized sums; this convention is necessary because the Fourier coefficients themselves are normalized.

The exact Fourier identity for the third uniformity norm is

$$\begin{aligned} \|f\|_{U^3}^8 &= \sum_{s,t \in \widehat{G}} |T(s,t)|^2, \\ T(s,t) &= \sum_{r \in \widehat{G}} a_r \overline{a_{r+s}} \overline{a_{r+t}} a_{r+s+t}. \end{aligned} \quad (\text{III.72.9})$$

For example, Fourier expansion of the original cube imposes four orthogonality equations. Solving these equations expresses its frequency variables as an affine cube; grouping its parallel faces gives (III.72.9). The solution uses additions and subtractions, so it introduces no torsion assumption.

On the two coordinate axes,

$$T(0,t) = T(t,0) = \sum_r \lambda_r \lambda_{r+t}.$$

If

$$E(\lambda) = \sum_{r+s=t+w} \lambda_r \lambda_s \lambda_t \lambda_w,$$

retaining those axes in the nonnegative sum gives

$$\|f\|_{U^3}^8 \geq 2E(\lambda) - S^2. \quad (\text{III.72.10})$$

The term S^2 corrects the double counting of $(s,t) = (0,0)$.

Theorem III.72.2 (A Fourier progression refinement). *Let f be real and centered on a finite abelian group of odd order, and define*

$$J(\lambda) = \sum_{r,t} \lambda_r^2 \lambda_t \lambda_{2r-t}.$$

Then

$$\|f\|_{U^3}^8 \geq 5S^2 - 10M + 4J(\lambda), \quad (\text{III.72.11})$$

$$\|f\|_{U^3}^8 \geq 5S^2 - 6M + 8 \sum_r |a_r|^6 |a_{3r}|^2. \quad (\text{III.72.12})$$

In particular, for any character ξ of odd order at least five, with $A = |a_\xi|$ and $B = |a_{3\xi}|$,

$$\|f\|_{U^3}^8 \geq 5S^2 - 6M + 16A^6 B^2. \quad (\text{III.72.13})$$

Proof. In $E(\lambda)$ retain the union of the three families

$$\mathcal{A} = \{r = t, s = w\}, \quad \mathcal{B} = \{r = w, s = t\}, \quad \mathcal{C} = \{s = -r, w = -t\}.$$

Each contributes S^2 , using the evenness of λ for $\mathcal{C}$. Every pairwise intersection contributes M . The triple intersection has $r = -r$, hence $r = 0$ because the group has odd order. Its weight is zero because $\lambda_0 = 0$. The union therefore has weight $3S^2 - 3M$.

Next retain every nonconstant repeated-entry quadruple

$$(r, r; t, 2r - t), \quad t \neq r, \quad (\text{III.72.14})$$

and its side reversal. Odd group order ensures that the two entries on the unequal side are distinct: $t = 2r - t$ would imply $t = r$. Thus these ordered quadruples are distinct from their side reversals, and have total weight $2(J - M)$. No positive-weight quadruple in this new family belongs to $\mathcal{A}$ or $\mathcal{B}$, since that would force $t = r$. Membership in $\mathcal{C}$ would force $r = -r$, hence $r = 0$ and zero weight. Consequently

$$E(\lambda) \geq 3S^2 - 5M + 2J(\lambda).$$

Combining this with (III.72.10) proves (III.72.11).

For each nonzero r , retain in J the diagonal $t = r$ and the additional indices $t = -r, 3r$. They are distinct: a coincidence requires $2r = 0$ or $4r = 0$. The $r = 0$ terms have zero weight. It follows that

$$J(\lambda) \geq M + 2 \sum_r \lambda_r^3 \lambda_{3r},$$

which proves (III.72.12). If $3r = 0$, the new summand vanishes, so the statement includes groups with three-torsion. Finally the terms $r = \xi, -\xi$ in that sum are equal and have value $A^6 B^2$, proving (III.72.13). $\square$

For a direct view of the last refinement, its extra energy configurations are obtained from

$$(\xi, \xi; -\xi, 3\xi), \quad (-\xi, -\xi; \xi, -3\xi)$$

by interchanging sides and ordering the unequal side. These are eight distinct configurations outside $\mathcal{A} \cup \mathcal{B} \cup \mathcal{C}$, each of weight $A^6 B^2$. The factor 16 in (III.72.13) includes the additional factor two from (III.72.10).

III.72.2 Two inequalities retaining the largest Fourier pair

Assume $f \neq 0$ and choose ξ maximizing $|a_\xi|$. Set

$$A = |a_\xi|, \quad p = \frac{2A^4}{S}, \quad D = 1 - p.$$

Here p is the fraction of Fourier fourth-power mass carried by the pair $\{\xi, -\xi\}$. In particular $0 < p \leq 1$. After multiplying f by a positive scalar we may suppose $S = 1$; then $A = (p/2)^{1/4}$.

If p_j are the corresponding masses of the other conjugate pairs, then $\sum_j p_j = D$ and

$$M = \frac{1}{2} \left(p^2 + \sum_j p_j^2 \right) \leq \frac{p^2 + D^2}{2}, \quad M \leq \frac{p}{2}. \quad (\text{III.72.15})$$

The second estimate uses the maximality of p . The energy argument without its new positive term gives

$$\|f\|_{U^3}^8 \geq 5 - 3p. \quad (\text{III.72.16})$$

If the order of ξ is at least five, (III.72.13) and the first inequality of (III.72.15) give

$$\|f\|_{U^3}^8 \geq 2 + 6pD + 16A^6 B^2, \quad B = |a_{3\xi}|. \quad (\text{III.72.17})$$

The quadratic dependence on B is the useful new feature. An estimate using only the residual fourth-power mass would permit B as large as $(D/2)^{1/4}$. Formula (III.72.17) imposes a further cost even when that mass is small.

III.72.3 The exact dual function of a cosine

For a real function g put

$$D_4 g(x) = \mathbb{E}_{h_1, h_2, h_3, h_4} \prod_{\substack{\omega \in \{0,1\}^4 \\ \omega \neq 0}} g(x + \omega \cdot h).$$

Cube Cauchy–Schwarz gives

$$|\mathbb{E}_x f(x) D_4 g(x)| \leq \|f\|_{U^4} \|g\|_{U^4}^{15}. \quad (\text{III.72.18})$$

The following finite identity was established in [162]. We give the counting proof because its coefficients are essential to the numerical improvement.

Lemma III.72.3 (Fourth-order cosine identity). *Let χ be a character of odd order at least five and let $\theta \in \mathbb{R}$. Write $g_j = e^{ij\theta} \chi^j + e^{-ij\theta} \chi^{-j}$. Then*

$$\|g_1\|_{U^4}^{16} = 222, \quad D_4 g_1 = 111g_1 + 27g_3 + g_5. \quad (\text{III.72.19})$$

When χ has order five, g_5 is constant and the formula still holds.

Proof. Expand the factors into characters. Each direction equation requires a sum of eight signs to be divisible by the odd order of χ . The sum is even and lies between -8 and 8 , so for an odd order at least five it must be zero as an integer. The counting is therefore unchanged for character orders five, seven, or larger.

For the full cube, each of two opposite three-faces has four plus signs. Let $n(t)$ count four-subsets of $\{0, 1\}^3$ with coordinate sum $t = (t_1, t_2, t_3)$. Splitting into these faces gives the total count $\sum_t n(t)^2$. The coefficient identity

$$[z^4] \prod_{\epsilon \in \{0,1\}^3} (1 + zx^\epsilon) = x_1 x_2 x_3 \prod_{j=1}^3 (1 + x_j)^2 + \sum_{i=1}^3 (1 + x_i^4) \prod_{j \neq i} x_j^2 \quad (\text{III.72.20})$$

gives the complete table

Type of t	Number of t	$n(t)$
One coordinate 0 or 4, the other two 2	6	1
All three coordinates in $\{1, 3\}$	8	1
Two in $\{1, 3\}$, one equal to 2	12	2
One in $\{1, 3\}$, two equal to 2	6	4
$(2, 2, 2)$	1	8

Thus the count is $6 + 8 + 12 \cdot 4 + 6 \cdot 16 + 64 = 222$. All surviving full-cube assignments have total frequency zero, which proves the first identity of (III.72.19).

For the dual function, adjoining the missing origin sign shows that the coefficients at frequencies 1 and -1 are each 111. At frequency -3 , the six plus-sign vertices among the fifteen nonzero vertices must place four selected vertices on each coordinate face. Complement all four coordinates. The resulting six vertices exclude 1111 and have coordinate sums $(2, 2, 2, 2)$. If a_j counts their vertices of Hamming weight j , the possible profiles and counts are

a_0	a_1	a_2	a_3	Count
0	4	2	0	3
1	2	3	0	12
1	3	1	1	12

In the first case the two edges form a perfect matching. In the second, choose two singleton endpoints in six ways and a three-edge path between them in two ways. In the third, choose the triple in four ways and the coordinate of that triple without its singleton in three ways; the remaining edge is forced. Hence the coefficient is $3 + 12 + 12 = 27$. Sign reversal gives the same coefficient at 3.

At frequency -5 , the five plus-sign vertices must have sixteen coordinate incidences. The unique possibility consists of the weight-four vertex and all four weight-three vertices. Its sign reversal gives frequency 5. Four distinct vertices have at most thirteen incidences, so no frequency of absolute value at least seven survives. This proves the dual identity. $\square$

Choose θ with $a_\xi = Ae^{i\theta}$ and apply the lemma to the corresponding character. Put $C = |a_{5\xi}|$. The pairs $\{\pm 3\xi\}$ and $\{\pm 5\xi\}$ are distinct and outside $\{\pm \xi\}$; the sole exception is $5\xi = 0$ at order five, in which case $C = 0$. The excluded coincidences would require $2\xi = 0$, $4\xi = 0$, $6\xi = 0$, or $8\xi = 0$. Consequently

$$2B^4 + 2C^4 \leq D. \quad (\text{III.72.21})$$

The dual pairing satisfies

$$\mathbb{E} f D_4 g_1 \geq 2(111A - 27B - C).$$

Therefore, with $R = \|f\|_{U^4}^{16}$ under the normalization $S = 1$, (III.72.18) gives

$$R \geq 222 \left[A - \frac{27}{111}B - \frac{1}{111}C \right]_+^{16}, \quad [x]_+ = \max(x, 0). \quad (\text{III.72.22})$$

III.72.4 The character of order three

The previous dual identity has a different form at order three. We treat that case directly, rather than imposing a stronger restriction on the ambient group.

Lemma III.72.4 (The ternary cosine). *For a character χ of order three and $g = e^{i\theta}\chi + e^{-i\theta}\chi^{-1}$,*

$$\begin{aligned} \|g\|_{U^4}^{16} &= 286 + 32 \cos(6\theta) + 16 \cos(12\theta) \\ &= 262 + 32 \left(\cos(6\theta) + \frac{1}{2} \right)^2 \geq 262. \end{aligned} \quad (\text{III.72.23})$$

If a largest Fourier pair of f has order three, then

$$R \geq \frac{131}{8}p^4. \quad (\text{III.72.24})$$

Proof. In the character expansion of the full cube, the possible numbers of plus signs are 2, 5, 8, 11, 14. Their multiplicities are respectively 8, 16, 286, 16, 8, as follows. At size two the vertices must be antipodal. At size five, reflect every coordinate that occurs four times; the result is the origin together with the four unit vertices, giving sixteen copies. At size eight, the 222 configurations from Lemma III.72.3 have four selected vertices on every coordinate face. Every remaining one consists of a whole three-face with a single vertex moved across the perpendicular direction. There are $4 \cdot 2 \cdot 8 = 64$ of these, each with exactly one exceptional direction. Taking complements gives sizes eleven and fourteen. The frequencies are $-12, -6, 0, 6, 12$, so summing their phases proves the first line of (III.72.23). The double-angle identity proves the second.

For (III.72.24), average f over the kernel of the chosen order-three character. Translation invariance and the triangle inequality make this averaging contractive for U^4 . On the three-element quotient the projected function is exactly Ag , because f is centered. Hence $R \geq 262A^{16} = (131/8)p^4$. $\square$

III.72.5 A scalar certificate using exact rational bounds

We now prove the constant 5. The proof reduces to two elementary one-variable estimates. Its few rational comparisons are also checked by the accompanying standard-library program `verification/verify_norm_gap.py`.

Proof of Theorem III.72.1. The zero function is immediate. Otherwise rescale to $S = 1$, so $R = \|f\|_{U^4}^{16}$. Suppose for a contradiction that $R \leq 5$. Monotonicity and the rational comparison $\sqrt{5} < 56/25$ give

$$\|f\|_{U^3}^8 \leq \|f\|_{U^4}^8 \leq \sqrt{5} < \frac{56}{25}.$$

Equation (III.72.16) implies $p > 23/25$. If the chosen character has order three, (III.72.24) yields

$$R > \frac{131}{8} \left(\frac{23}{25} \right)^4 > 5,$$

a contradiction. Its order is therefore at least five.

Since $p > 1/2$ and $p(1-p)$ decreases on $[1/2, 1]$, (III.72.17) implies $p > 19/20$. Indeed, $p \leq 19/20$ would give $6p(1-p) \geq 57/200 > 6/25$, contradicting that equation. Thus

$$0 \leq D < \frac{1}{20}, \quad B^2 < \frac{6/25 - (57/10)D}{16A^6}. \quad (\text{III.72.25})$$

In particular $D < 4/95$.

Put $a = 2^{-1/4}$. For $0 \leq D \leq 1/20$,

$$A = a(1-D)^{1/4} \geq a(1-D/3). \quad (\text{III.72.26})$$

To check this, differentiate $(1-D)^{1/4} - (1-D/3)$. Its derivative is positive on this interval because $(1-D)^{3/4} \geq 19/20 > 3/4$. We shall prove

$$A - \frac{27}{111}B - \frac{1}{111}C > \frac{789}{1000}. \quad (\text{III.72.27})$$

This contradicts (III.72.22) and $R \leq 5$, since

$$222 \left(\frac{789}{1000} \right)^{16} > 5. \quad (\text{III.72.28})$$

First suppose $0 \leq D \leq 1/300$. Hölder's inequality and (III.72.21) give

$$27B + C \leq (27^{4/3} + 1)^{3/4}(B^4 + C^4)^{1/4} \leq 82^{3/4}(D/2)^{1/4}.$$

The exact comparisons

$$a > \frac{21}{25}, \quad 82^{3/4} < \frac{273}{10}, \quad (1/600)^{1/4} < \frac{203}{1000}$$

and (III.72.26) show that the left side of (III.72.27) is greater than

$$\frac{21}{25} \frac{899}{900} - \frac{1}{111} \frac{273}{10} \frac{203}{1000} = \frac{18879}{22500} - \frac{55419}{1110000} > \frac{789}{1000}. \quad (\text{III.72.29})$$

It remains to consider $1/300 \leq D < 4/95$. Equations (III.72.25), (III.72.26), and $C \leq (D/2)^{1/4}$ show that the left side of (III.72.27) is at least

$$\begin{aligned} F(D) &= a(1-D/3) - \frac{27}{444a^3}(1-D/3)^{-3}\sqrt{6/25 - 57D/10} \\ &\quad - \frac{1}{111}(D/2)^{1/4}. \end{aligned} \quad (\text{III.72.30})$$

For $h(D) = (1-D/3)^{-3}\sqrt{6/25 - 57D/10}$, differentiation gives

$$-h'(D) = (1-D/3)^{-4} \frac{261/100 + (19/4)D}{\sqrt{6/25 - 57D/10}} > 5.$$

The last inequality follows from $\sqrt{6/25 - 57D/10} < 1/2$ and $(1-D/3)^{-4} \geq 1$. Use $a^3 < 3/5$, $a < 1$, and $300^{3/4} < 73$ to obtain

$$F'(D) > \frac{225}{444} - \frac{1}{3} - \frac{73}{444} = \frac{1}{111} > 0. \quad (\text{III.72.31})$$

Here the last negative term bounds the derivative of $(D/2)^{1/4}/111 = aD^{1/4}/111$ for $D \geq 1/300$.

It suffices to evaluate F at $D = 1/300$. The exact bounds

$$\begin{aligned} a \frac{899}{900} &> \frac{4199}{5000}, \\ \frac{\sqrt{221/1000}}{4a^3(899/900)^3} &< \frac{1}{5}, \\ (1/600)^{1/4} &< \frac{203}{1000} \end{aligned}$$

give

$$F(1/300) > \frac{4199}{5000} - \frac{27/5 + 203/1000}{111} = \frac{4199}{5000} - \frac{5603}{111000} > \frac{789}{1000}. \quad (\text{III.72.32})$$

For the first bound use $a > 1051/1250$. For the second, squaring reduces it to

$$\frac{221}{1000} < \frac{16}{25} a^6 \left(\frac{899}{900} \right)^6,$$

which holds already with a replaced by $21/25$. Thus every comparison in the proof reduces to integer arithmetic after clearing positive denominators and, where necessary, taking fourth powers.

Equations (III.72.31) and (III.72.32) finish (III.72.27), contradicting $R \leq 5$. Undoing the normalization proves $\|f\|_{U^4}^{16} > 5S^4$ for $f \neq 0$. Taking fourth roots yields $\|f\|_{U^4}^4 \geq 5^{1/4}S$. Finally $\|f\|_{U^d} \geq \|f\|_{U^4}$ for $d \geq 4$ proves (III.72.4) with exactly the asserted exponents. $\square$

III.72.6 The cube-count consequence and the remaining extremal problem

For a real function F on G , write

$$Q_d(F) = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{\omega \in \{0,1\}^d} F(x + \omega \cdot h).$$

The fourth-order expansion proved in [159] says that, on odd-order groups, for $f = F - \delta$, $\mathbb{E}f = 0$, $m = 2^d$, and $u = \|f\|_{U^d}$,

$$\begin{aligned} Q_d(F) &= \delta^m + R_d \delta^{m-4} \|f\|_{U^2}^4 + \mathcal{R}_d, \\ |\mathcal{R}_d| &\leq \sum_{j=5}^m \binom{m}{j} \delta^{m-j} u^j, \quad R_d = \frac{6^d - 2 \cdot 4^d + 2^d}{8}. \end{aligned} \quad (\text{III.72.33})$$

This statement assumes $\delta \geq 0$.

Corollary III.72.5 (An improved leading cube-count bound). *Let $d \geq 4$, let G have odd order, and let $F : G \rightarrow \mathbb{R}$ have mean $\delta \geq 0$. With $f = F - \delta$, $m = 2^d$, and $u = \|f\|_{U^d}$,*

$$Q_d(F) \leq \delta^m + 5^{-1/4} R_d \delta^{m-4} u^4 + \sum_{j=5}^m \binom{m}{j} \delta^{m-j} u^j. \quad (\text{III.72.34})$$

In particular, at $d = 4$ the quartic multiplier is $100 \cdot 5^{-1/4} = 66.8740304976\dots$, improving the strongest inspected prior multiplier $100(8/35)^{1/4} = 69.1441569\dots$ from [163]. For $d \geq 5$, retain instead the smaller multiplier supplied by the minimum in (III.72.8).

Proof. Substitute Theorem III.72.1 into the exact quartic term of (III.72.33), leaving the established remainder unchanged. $\square$

[write] **6 October 2026; sources 34–39.** Source 34’s Corollary III.71.9 has $200/3 = 66.667$ in place of $100 \cdot 5^{-1/4} = 66.874$ here (recomputed at the write).

For a set $A \subseteq G$ satisfying $\|\mathbf{1}_A - \delta\|_{U^d}^m \leq \alpha$, replace u^j in (III.72.34) by $\alpha^{j/m}$. Multiplication by $|G|^{d+1}$ converts the normalized average into the parameterized cube count used in the paper. This is a refinement of the local Lemma 3.10 interface. It neither changes the order of the quartic error nor supplies an optimal coefficient for $d \geq 4$.

Methods and limitations. The new input is the explicit retention of repeated-entry additive quadruples in Theorem III.72.2. The cosine dual identity and the preceding norm-gap framework are credited to the earlier reports and reproved where used. All constants in the final scalar argument are certified by exact arithmetic; the program also enumerates the sign counts behind Lemmas III.72.3 and III.72.4. The proof is mathematical rather than Lean-formalized. It provides a new bound relative to the inspected repository materials, with no claim of an exhaustive literature-priority determination.

Research question III.72.6 (The next quantitative and structural questions). Can the full statistic $J(\lambda)$, or further nonnegative terms in (III.72.9), determine the optimal c_4^{odd} ? Does an extremizer live on a finite cyclic quotient, or only arise as a limit? At higher orders, can the additional harmonics in D_d of a cosine yield constants that improve with d , rather than the single fourth-order constant propagated by monotonicity? The negative third-harmonic examples in [162] show that any proposed classification must retain Fourier phases as well as amplitude masses.

III.73 Source 34’s chapter: prime torsion in cube expansions [write]

Added in the seventh write, 6 October 2026.

The next eight sections are source 34’s Sections 2–9, on when a prime p first enters the centered expansion of a cube count through Smith torsion of a Boolean configuration. *Logically the chapter follows source 09’s support-lattice material* (Sections III.18 and III.23: integer value lattices, the order-five obstruction and the all-prime family with $h(\ell) \leq \ell + 2$) and source 21’s order-five certificates; it is printed here because no number printed earlier may move. Source 34’s introduction is Section F.39; its Section 10 is in the chapter on c_d^{odd} (Section III.71).

What is proved, in one paragraph. The least support $h(p)$ of an affinely independent Boolean configuration admitting a full-support characteristic- p relation is $D_{\text{div}}(p)+1$, where $D_{\text{div}}(p)$ is the least order of a nonsingular binary matrix with determinant divisible by p (Theorem III.75.2); $h(7) = 7$ and $h(11) = h(13) = h(17) = 8$, with exhaustive certificates for 11 and 13 (Theorem III.76.2); every integer up to $\exp(\frac{1}{2}N \log N - N \log \log N - 4N)$ is eventually the sole nonunit Smith factor of a binary matrix of order N , through Shitov’s composition and the Alon–Vu matrices (Theorem III.78.1); hence $h(p) \sim h_{\text{cyc}}(p) \sim 2 \log p / \log \log p$ (Theorem III.79.1); universal nonnegativity of cyclic moments holds exactly for inverse-paired coefficients, and $h_{\text{pos}}(m) \sim 4 \log m / \log \log m$ (Theorems III.80.1, III.80.3); the exact quartic term is recalled with proof (Proposition III.81.2). This answers source 09’s prime-onset question (item III.16).

Relation to the formal statements (checked at HEAD 89c1cbc38). lemma_3_10 (Sections01_03.

lean, line 340) has the exact companion lemma_3_10_holds (Proofs03CubeUpper.lean, line 82); no declaration concerns Smith torsion of cube expansions. Nothing in this chapter is formalized.

III.74 Fourier constraints and the Boolean row quotient

Source 34, delivered Section 2.

Throughout, G is a finite abelian group, written additively, and $\widehat{G}$ is its character group, written multiplicatively. Every average is normalized. For $f : G \rightarrow \mathbb{C}$ set

$$\widehat{f}(\chi) = \mathbb{E}_{x \in G} f(x) \overline{\chi(x)}.$$

The trivial character is 1. Thus $\widehat{f}(1) = 0$ for centered f , and $\widehat{f}(\chi^{-1}) = \overline{\widehat{f}(\chi)}$ for real f . The Gowers norm is normalized by

$$\|f\|_{U^d}^{2^d} = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{v \in \{0,1\}^d} \mathcal{C}^{|v|} f(x + v \cdot h),$$

where $\mathcal{C}$ denotes complex conjugation. In particular,

$$\|f\|_{U^2}^4 = \sum_{\chi \in \widehat{G}} |\widehat{f}(\chi)|^4.$$

For real functions the conjugations have no effect. We use the standard multilinear Gowers–Cauchy–Schwarz inequality and norm monotonicity; these follow by successive Cauchy–Schwarz applications as in Gowers’s Lemmas 3.8–3.9 [1].

Definition III.74.1 (Augmented Boolean matrix). Let $S = \{v_1, \dots, v_s\} \subseteq \{0, 1\}^d$, with distinct vertices. Its augmented matrix and row quotient are

$$M_S = \begin{pmatrix} 1 & \cdots & 1 \\ v_1 & \cdots & v_s \end{pmatrix}, \quad L_S = \text{row}_{\mathbb{Z}}(M_S) \subseteq \mathbb{Z}^s, \quad A_S = \mathbb{Z}^s / L_S.$$

The columns are written vertically. The set S is affinely independent exactly when $\text{rank}_{\mathbb{Q}} M_S = s$.

Proposition III.74.2 (Exact Fourier presentation). *For any functions $f_1, \dots, f_s : G \rightarrow \mathbb{C}$,*

$$\mathbb{E}_{x, h} \prod_{j=1}^s f_j(x + v_j \cdot h) = \sum_{\substack{\chi_1, \dots, \chi_s \in \widehat{G} \\ M_S \chi = 1}} \prod_{j=1}^s \widehat{f}_j(\chi_j). \quad (\text{III.74.1})$$

Here $M_S \chi = 1$ means that the product of the characters in every row, with the indicated exponents, is trivial. The solution tuples are canonically $\text{Hom}(A_S, \widehat{G})$, evaluated at the images of the standard generators of $\mathbb{Z}^s$. If all f_j are centered, only tuples with every $\chi_j \neq 1$ contribute.

Proof. Insert Fourier inversion into each factor. Averaging first in x forces $\prod_j \chi_j = 1$. Averaging in h_i forces $\prod_j \chi_j^{(v_j)_i} = 1$. These are exactly the rows of M_S . A tuple defines a homomorphism from $\mathbb{Z}^s$ to $\widehat{G}$; the row constraints say it factors through A_S . Centering removes any summand with a trivial coordinate character. $\square$

Definition III.74.3 (Minimal supports). For a prime p , let $h(p)$ be the least s for which an affinely independent Boolean set S of size s has a vector $0 \neq c \in \ker_{\mathbb{F}_p} M_S$ with all coordinates nonzero. Equivalently, A_S has p -torsion and a full-support characteristic- p character. Define

$$D_{\text{div}}(p) = \min\{n : B \in \{0,1\}^{n \times n}, \det B \neq 0, p \mid \det B\}.$$

For an integer $m \geq 2$, let $h_{\text{cyc}}(m)$ be the least s for an affinely independent Boolean S such that $A_S \cong \mathbb{Z}/m\mathbb{Z}$ and the images of all standard coordinate generators are nonzero in this quotient. In terms of a faithful character to $\mathbb{Z}/m\mathbb{Z}$, its coefficient vector has no zero coordinate. For composite m , these coordinates need not be units. For primes, $h(p) \leq h_{\text{cyc}}(p)$.

These minima exist. For example, the zero vertex and the columns of $J - I_{p+1}$ give the earlier $p+2$ construction. Section III.77 gives a construction for every cyclic order. All logarithms below are natural unless a base is displayed.

Remark III.74.4 (Which characteristic dependence is measured). The parameter $h(p)$ measures *Smith torsion*: a drop of the rank of the integer constraint matrix modulo p . A torsion-free integer relation can also have coefficients that vanish modulo p , changing whether a centered Fourier term survives. Our dimension threshold does not rule out every possible characteristic-dependent phenomenon of that latter kind. This distinction is essential when using an all-orders expansion.

III.75 The exact determinant reduction

Source 34, delivered Section 3.

The central observation is that minimal order removes the extra support obstruction. A general matrix singular modulo p can have only relations with zero coordinates; a dimension-minimal binary matrix with determinant divisible by p cannot.

Lemma III.75.1 (Removing the affine row). *Suppose a square, nonsingular binary matrix M has order s and its first row is all ones. There is a binary matrix B of order $s-1$ with $|\det B| = |\det M|$.*

Proof. For each lower row whose first entry is one, replace that row by the first row minus itself. All entries remain binary and the absolute determinant is unchanged. The first column is now $(1, 0, \dots, 0)^T$. Subtract it from every other column and expand along the first row. The remaining lower-right matrix is binary and has the claimed determinant. $\square$

Theorem III.75.2 (Full support is automatic at minimal order). *For every prime p ,*

$$h(p) = D_{\text{div}}(p) + 1. \tag{III.75.1}$$

If B attains $n = D_{\text{div}}(p)$, then $\ker_{\mathbb{F}_p} B$ is one-dimensional, every nonzero vector in it has full support, and its coordinate sum is nonzero. The zero vertex and the columns of B give a witness for $h(p)$ in dimension n .

Proof. Take an admissible set of s vertices. Since the augmented matrix has rank s over $\mathbb{Q}$, the nonzero all-ones row can be extended to a basis of its row space. The resulting $s \times s$ minor is nonsingular over $\mathbb{Q}$ and singular over $\mathbb{F}_p$, because it annihilates the given nonzero vector. Lemma III.75.1 produces a nonsingular binary matrix of order $s-1$ with determinant divisible by p . Thus $D_{\text{div}}(p) \leq s-1$.

Conversely, take B of minimal order n and $0 \neq c \in \ker_{\mathbb{F}_p} B$. If c has support of size $r < n$, the corresponding r columns of B are linearly independent over $\mathbb{Q}$. Select r rows on which those columns give a nonsingular $r \times r$ submatrix. Its reduction modulo p annihilates the nonzero restricted vector. This contradicts minimality. Hence every nonzero c has full support.

Suppose now that $\sum_i c_i = 0$. Adjoin an all-ones row to B . The new row space still has rank n over $\mathbb{Q}$, and we may choose n independent rows including the all-ones row. They form a square binary matrix singular modulo p . Applying Lemma III.75.1 again gives a forbidden matrix of order $n - 1$. Therefore $\sum_i c_i \neq 0$.

The augmented matrix of the zero vertex and the columns of B is

$$M = \begin{pmatrix} 1 & \mathbf{1}^T \\ 0 & B \end{pmatrix}.$$

It is nonsingular over $\mathbb{Q}$ and annihilates $(-\sum_i c_i, c_1, \dots, c_n)$ modulo p . This vector has full support, giving $h(p) \leq n + 1$.

Finally, if the kernel of B had dimension at least two, two independent kernel vectors would have a nonzero linear combination whose first coordinate is zero. That contradicts the full-support conclusion just proved. The nullity is therefore one. $\square$

Corollary III.75.3 (Arbitrary support and the first cube dimension). *Let $S \subseteq \{0, 1\}^d$ be arbitrary, possibly affinely dependent, and let $r = \text{rank}_{\mathbb{Q}} M_S$. If A_S has nonzero p -primary torsion, then*

$$r \geq D_{\text{div}}(p) + 1 = h(p), \quad d \geq D_{\text{div}}(p), \quad |S| \geq h(p). \quad (\text{III.75.2})$$

All three minimum values are attained simultaneously by the construction in Theorem III.75.2.

Proof. Smith normal form says that A_S has p -torsion exactly when $\text{rank}_{\mathbb{F}_p} M_S < r$. Choose r independent rational rows including the all-ones row, and then r columns on which their minor is nonsingular. Every $r \times r$ minor vanishes modulo p . Lemma III.75.1 yields a nonsingular binary determinant of order $r - 1$ divisible by p . Hence $r - 1 \geq D_{\text{div}}(p)$. The other inequalities follow from $r \leq d + 1$ and $r \leq |S|$. The theorem provides equality with $r = |S| = d + 1$. $\square$

Remark III.75.4. This corollary concerns *any* p -primary Smith torsion, so introducing extra vertices or rational affine dependencies cannot make it occur in a smaller cube dimension. It does not say that a minimum determinant divisible by p has absolute value p ; other prime factors or higher powers may remain. Exact cyclic isolation is treated separately.

Lemma III.75.5 (Cyclic support compression). *Let B be a nonsingular binary matrix of order n with $\mathbb{Z}^n / \text{row}_{\mathbb{Z}}(B) \cong \mathbb{Z}/m\mathbb{Z}$. Then*

$$h_{\text{cyc}}(m) \leq n + 1. \quad (\text{III.75.3})$$

More precisely, starting from the zero vertex and the columns of B , delete the coordinates on which a faithful quotient character is zero. The retained Boolean set is affinely independent, has quotient exactly $\mathbb{Z}/m\mathbb{Z}$, and has a full-support faithful character.

Proof. The augmented matrix M is integrally equivalent to $\text{diag}(1, B)$, so its row quotient is cyclic of order m . Choose an isomorphism to $\mathbb{Z}/m\mathbb{Z}$ and represent the resulting character on $\mathbb{Z}^{n+1}$ by a vector c modulo m . Let T be its nonzero coordinates and let π be projection onto them. Writing $L = \text{row}_{\mathbb{Z}}(M)$, there is a surjection

$$\mathbb{Z}^{n+1}/L \longrightarrow \mathbb{Z}^T/\pi(L).$$

The isomorphism to $\mathbb{Z}/m\mathbb{Z}$ factors through this surjection, because c is zero outside T . Thus the target is both a quotient of a group of order m and a group surjecting onto one of order m . Both maps must be isomorphisms. Projection of the rows is exactly deletion of the unwanted columns of M . The corresponding subset of an affinely independent set remains affinely independent, and the retained vector has full support by construction. $\square$

III.76 Exact first cases and finite nonexistence certificates

Source 34, delivered Section 4.

Lemma III.76.1 (Binary Hadamard bound). *Every binary matrix B of order n satisfies*

$$|\det B| \leq \frac{(n+1)^{(n+1)/2}}{2^n}.$$

Consequently every admissible support of size s for $h(p)$, or for $h_{\text{cyc}}(m)$, satisfies respectively

$$p \leq \frac{s^{s/2}}{2^{s-1}}, \quad m \leq \frac{s^{s/2}}{2^{s-1}}. \quad (\text{III.76.1})$$

Proof. The sign matrix

$$H = \begin{pmatrix} 1 & \mathbf{1}^T \\ \mathbf{1} & J - 2B \end{pmatrix}$$

has $|\det H| = 2^n |\det B|$. Every row has Euclidean norm $\sqrt{n+1}$, so Hadamard's inequality gives the first assertion. For $h(p)$ use the nonsingular minor of Theorem III.75.2. For $h_{\text{cyc}}(m)$, each full minor of the augmented matrix is divisible by its full row-lattice index m : this follows from Smith normal form, or from the gcd-of-minors formula for the index. Choose a nonzero full minor including the all-ones row and apply Lemma III.75.1. $\square$

Theorem III.76.2 (Four new exact prime values). *One has*

$$h(7) = h_{\text{cyc}}(7) = 7, \quad h(11) = h_{\text{cyc}}(11) = h(13) = h_{\text{cyc}}(13) = h(17) = h_{\text{cyc}}(17) = 8.$$

Together with the earlier values, the initial table is

	p	2	3	5	7	11	13	17
$h(p) = h_{\text{cyc}}(p)$		4	5	6	7	8	8	8
<i>minimum cube dimension</i>		3	4	5	6	7	7	7

Proof. The upper bounds follow from the explicit matrices below. They have determinants 7, 11, 13, 17, respectively, and therefore cyclic row quotients of those prime orders. Their augmented coefficient vectors have full support, so no support compression is needed.

For the lower bounds, the Hadamard bound at $s = 6$ is $27/4 < 7$, and at $s = 7$ is $7^{7/2}/64 < 17$. The same bound increases with $s \geq 2$, so it also excludes smaller supports. For $p = 11, 13$, the classical binary maximum determinant of order six is nine [145]; Theorem III.75.2 then excludes $s \leq 7$. A fully independent finite proof of that exclusion, specialized to these two primes, is supplied immediately below and in the certificate files. The values at 2, 3, 5 and their cyclic witnesses were established in source 09 of [155]. $\square$

[write] **6 October 2026; sources 34–39.** *Answers source 09’s question* (its Research question “Support-efficient torsion”, item III.16): $h(7) = 7$ and $h(11) = 8$, and Theorem III.79.1 gives the leading asymptotic of the gap between source 09’s lower bound (III.23.4) and $h(\ell) \leq \ell + 2$. The intake recomputed, in exact integer arithmetic, the four printed matrices (determinant p , full-support kernel modulo p , cyclic quotient) and the orders-1–4 binary maximal determinants; the exhaustive certificates for $p = 11, 13$ are checked by the source’s own verifier (rerun on a copy at the write). Still open: whether $h(p) = h_{\text{cyc}}(p)$ for every prime, and the second-order term (Research questions V.130.1 and V.130.2).

III.76.1 Positive matrices and full-support relations

For each displayed B_p , take the zero vertex and its columns. Let $M_p = \begin{pmatrix} 1 & \mathbf{1}^T \\ 0 & B_p \end{pmatrix}$. The matrices are

$$B_7 = \begin{pmatrix} 0 & 0 & 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 & 0 \end{pmatrix}, \quad B_{11} = \begin{pmatrix} 0 & 0 & 0 & 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 1 & 0 & 0 & 1 \\ 0 & 1 & 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 & 0 & 0 & 1 \\ 1 & 1 & 1 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 0 \end{pmatrix},$$

$$B_{13} = \begin{pmatrix} 0 & 0 & 0 & 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 1 & 0 & 0 & 1 \\ 1 & 1 & 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 0 \end{pmatrix}, \quad B_{17} = \begin{pmatrix} 1 & 0 & 0 & 0 & 0 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 0 \end{pmatrix}.$$

The coefficient vectors, starting with the zero vertex, are

$$\begin{aligned} c_7 &= (1, 1, 1, 1, 3, 3, 4), \\ c_{11} &= (1, 1, 1, 1, 3, 4, 7), \\ c_{13} &= (1, 1, 1, 2, 3, 5, 8), \\ c_{17} &= (1, 1, 2, 4, 4, 6, 7, 9). \end{aligned}$$

In each case $M_p c_p = (2p, p, \dots, p)^T$. In particular, every coordinate is nonzero modulo p . Exact fraction-free determinants are recorded in `data/small_prime_matrices.json` and checked by `code/verify_torsion.py`.

The first three determinant computations also have short elimination proofs. In $B_7 x = 0$, set $x_6 = t$. The first two rows give $x_4 = x_5 = -t$, the last three give $x_1 = x_2 = x_3 = 2t$, and the remaining row gives $7t = 0$. In $B_{11} x = 0$, setting $x_7 = t$ gives $x_5 = x_6 = -t$, $x_4 = 2t$, and $x_1 = x_2 = x_3 = -3t$, with final relation $11t = 0$. For B_{13} the corresponding values are $x_5 = x_6 = -t$, $x_4 = 2t$, $x_3 = -3t$, and $x_1 = x_2 = 5t$, with final relation $13t = 0$. These are eliminations over $\mathbb{Z}$, with pivots of absolute value one, so they identify the exact cyclic quotient rather than only a modular kernel. The fourth matrix is verified by the same exact integer determinant algorithm included in the package.

III.76.2 A small certificate for nonexistence at seven vertices

Here is a direct proof method for the $p = 11, 13$ lower bounds that does not rely on a maximum-determinant table or floating-point rank. Suppose M is an augmented seven-column Boolean matrix with full rank over $\mathbb{Q}$ and a full-support kernel vector modulo p . The all-ones row forces

$$c \in (\mathbb{F}_p^\times)^7, \quad \sum_j c_j = 0.$$

Permuting coordinates and multiplying all c_j by a nonzero scalar do not change the existence of M . Represent these equivalence classes by the lexicographically least sorted tuple in each scalar orbit. There are exactly 104 such tuples for $p = 11$ and 204 for $p = 13$.

For a representative c , let

$$Z_p(c) = \{a \in \{0, 1\}^7 : a \cdot c = 0 \pmod{p}\}.$$

Every possible row of M belongs to this explicitly finite set. The supplied certificate gives, for each representative, a nonzero integer vector $z_c \in \mathbb{Z}^7$ satisfying

$$a \cdot z_c = 0 \quad \text{over } \mathbb{Z} \quad \text{for every } a \in Z_p(c). \quad (\text{III.76.2})$$

Thus all rows of M lie in a proper rational hyperplane, a contradiction.

The verifier independently regenerates the complete list of scalar and permutation orbit representatives. For each, it checks every one of the 2^7 Boolean masks and the integer dot product in (III.76.2). It verifies that z_c is a nonzero integer vector. It does *not* use a numerical rank routine, nor does it trust the optional rank field stored by the certificate generator. Consequently the negative proof consists of a short exhaustiveness argument and finitely many exact integer identities. The files are `data/no_seven_vertex_torsion_mod_11.json` and `data/no_seven_vertex_torsion_mod_13.json`.

III.76.3 What these configurations contribute to a cube expansion

Because $A_{S_p} \cong \mathbb{Z}/p\mathbb{Z}$ with generator coordinates c_p , Proposition III.74.2 gives, for a real centered f ,

$$T_{S_p}(f) := \mathbb{E}_{x,h} \prod_{v \in S_p} f(x + v \cdot h) = \sum_{\substack{\chi \in \widehat{G} \\ \chi^p = 1}} \prod_{j=1}^{|S_p|} \widehat{f}(\chi^{(c_p)_j}). \quad (\text{III.76.3})$$

For example, on $G = \mathbb{Z}/p\mathbb{Z}$, the centered real function $f = p\mathbf{1}_{\{0\}} - 1$ has Fourier coefficient one at every nontrivial character. The value of (III.76.3) is therefore $p - 1$. On a general finite group the analogous function $|G|\mathbf{1}_{\{0\}} - 1$ gives $|\widehat{G}[p]| - 1 = |G[p]| - 1$.

These are also legitimate perturbative examples for bounded functions. On $\mathbb{Z}/p\mathbb{Z}$, if $0 < \delta < 1$ and $0 < t \leq \min\{\delta, (1 - \delta)/(p - 1)\}$, then $F = \delta + tf$ takes values in $[0, 1]$, and the selected support contributes a nonzero term of degree $|S_p|$ in t . In a heterogeneous cube product, placing $\delta_v + t_v f$ at selected vertices makes the corresponding mixed coefficient individually visible. This avoids any claim that different supports cannot cancel in an arbitrary homogeneous expansion.

III.77 An elementary positive detector for every cyclic order

Source 34, delivered Section 5.

The asymptotically optimal construction later in the article uses Alon–Vu matrices. The following elementary construction is useful when a short explicit presentation and positivity matter more than the asymptotic number of vertices.

Theorem III.77.1 (Positive cyclic detector). *For every integer $q \geq 2$, put $r = \lceil \log_2 q \rceil$ and $t = r + 1$. There is an affinely independent Boolean set S_q of $2t$ vertices in dimension $2t - 1$ such that $A_{S_q} \cong \mathbb{Z}/q\mathbb{Z}$, with full-support generator coordinates*

$$(a_1, -a_1, \dots, a_t, -a_t), \quad (a_1, \dots, a_t) = (1, 1, 2, 4, \dots, 2^{r-1}).$$

For every real centered $f : G \rightarrow \mathbb{R}$ its moment is

$$T_{S_q}(f) = \sum_{\substack{\chi \in \widehat{G} \\ \chi^q = 1}} \prod_{i=1}^t |\widehat{f}(\chi^{a_i})|^2 \geq 0. \quad (\text{III.77.1})$$

In particular,

$$h_{\text{cyc}}(q) \leq 2\lceil \log_2 q \rceil + 2.$$

Proof. Use $2t$ generators x_i, y_i and the following binary relations:

$$x_i + y_i = 0 \quad (1 \leq i \leq t), \quad (\text{III.77.2})$$

$$\sum_{j < i} x_j + y_i = 0 \quad (2 \leq i \leq t), \quad (\text{III.77.3})$$

$$\sum_{i \in J} x_i = 0. \quad (\text{III.77.4})$$

Choose J so that $\sum_{i \in J} a_i = q$. Such a subset exists: the multiset $1, 1, 2, \dots, 2^{r-1}$ has every subset-sum value from zero to 2^r , including q .

Relations (III.77.2) give $y_i = -x_i$. Relations (III.77.3) give $x_i = \sum_{j < i} x_j = a_i x_1$. The last relation is exactly $q x_1 = 0$. Each elimination has an integer unit pivot, so the square binary relation matrix has determinant of absolute value q and row quotient exactly $\mathbb{Z}/q\mathbb{Z}$.

Replace its first pair row by the sum of all pair rows. This is a unimodular row operation, and the replacement row is all ones. Treat the remaining $2t - 1$ binary rows as coordinates of vertices. The square augmented matrix is nonsingular over $\mathbb{Q}$, so its columns are distinct and give an affinely independent Boolean set. Every a_i is positive and less than q , including when q is a power of two. Thus the quotient character has full support.

Finally, apply Proposition III.74.2 and use reality of f to pair the coefficients at χ^{a_i} and χ^{-a_i} . This proves (III.77.1). $\square$

Remark III.77.2 (Odd and even cyclic orders). For odd q , every a_i is a unit modulo q . On $\mathbb{Z}/q\mathbb{Z}$ the function $f = q\mathbf{1}_{\{0\}} - 1$ therefore gives $T_{S_q}(f) = q - 1$. For even q , a nontrivial character of a proper divisor order can be killed by some a_i ; the value need not be $q - 1$. Every character of exact order q still contributes one, so the detector is nonzero and at least $\varphi(q)$ for this function. This distinction prevents an unwarranted prime-to-composite extension of the sharper positivity statement.

The construction supplies an explicit all-parameter proof. The accompanying program checks it for 36 moduli, including 65537, using exact determinants and modular syndromes. It also computes the complete spatial average directly at $q = 2, 3$, independently of Fourier inversion. These checks validate the implementation; the integer elimination proves the theorem for every q .

III.78 Consecutive binary determinants with controlled Smith form

Source 34, delivered Section 6.

The determinant characterization reduces the support problem to a question about binary matrices. For the asymptotic upper bound, a construction that produces every sufficiently small integer is more useful than a construction producing only one exceptionally large determinant. Shah’s published consecutive range is of order $2^N/N$ [146]. Shitov’s author-uploaded 2025 preprint obtains the superexponential range

$$\exp\left(\frac{N}{2} \log N - 7N \log \log N\right)$$

for all sufficiently large orders N [147]. Its construction uses the unimodular matrices with large inverse entries of Alon and Vü [144].

We retain that construction, give all integer-rounding and determinant sign details, and improve its stated secondary loss. Tracking the integer row and column operations gives a further conclusion about the whole Smith normal form. The sole deep external input in the following proof is the published Alon–Vü theorem; the remaining arguments are included in full.

Theorem III.78.1 (Consecutive determinants and cyclic Smith forms). *For every sufficiently large integer N , and every integer h satisfying*

$$1 \leq h \leq \left\lfloor \exp\left(\frac{N}{2} \log N - N \log \log N - 4N\right) \right\rfloor, \quad (\text{III.78.1})$$

there exists $B \in \{0, 1\}^{N \times N}$ with

$$\det B = h, \quad \text{SNF}(B) = \text{diag}(1, \dots, 1, h).$$

The matrix may moreover be chosen to have an $(N-1) \times (N-1)$ minor equal to 1. Consequently its integer row-lattice quotient is cyclic of order h . The determinant conclusion also holds for $h = 0$.

The theorem improves the secondary term in the displayed range; its leading term $\frac{1}{2}N \log N$ was already obtained by Shitov. Cyclicity for composite h is additional information: the equality $\det B = h$ alone does not imply a cyclic quotient.

III.78.1 A unimodular matrix supplies a large integer

We use the arbitrary-order construction in Section 2.4 of Alon–Vü, together with the remark at the end of that section [144]. In natural logarithms, its relevant consequence is the existence, for every sufficiently large m , of a binary matrix M with $|\det M| = 1$ and

$$\max_{i,j} |(M^{-1})_{ij}| \geq \exp\left(\frac{m}{2} \log m - (2 \log 2 + o(1))m\right). \quad (\text{III.78.2})$$

The unimodularity assertion matters: it permits later substitutions by integer row and column operations.

Lemma III.78.2 (A large integer with a unimodular auxiliary block). *For every sufficiently large integer q , there are a binary $(q-1) \times (q-1)$ matrix M with $\det M = 1$ and binary unit column vectors u, v such that, on putting*

$$\alpha = -u^\top M^{-1}v, \quad A = |\alpha|,$$

one has $A > 1$ and

$$\frac{q}{2} \log q - \frac{3}{2}q \leq \log A \leq \frac{q}{2} \log q. \quad (\text{III.78.3})$$

Proof. Apply (III.78.2) with $m = q - 1$. If necessary, interchange two rows to arrange $\det M = 1$. This operation permutes the absolute values of the inverse entries, so it preserves the required lower bound. Choose u, v to select a suitable inverse entry. The binary matrix

$$U = \begin{pmatrix} 0 & u^\top \\ v & M \end{pmatrix}$$

has determinant α , by elimination using the invertible block M . Since $2 \log 2 < 3/2$, replacing $m = q - 1$ by q in the lower estimate costs only $O(\log q)$ and gives the left side of (III.78.3) for all sufficiently large q . The row-norm form of Hadamard's inequality applied to U gives $A \leq q^{q/2}$, proving the upper bound. The lower bound tends to infinity, so $A > 1$ eventually. $\square$

The integer α may have either sign. The construction below handles both signs without changing any binary entries.

III.78.2 Encoding one bounded row by binary rows

Lemma III.78.3 (Eliminating small negative entries). *Let R be a square integer matrix with a selected row z whose entries belong to $\{-2, -1, 0, 1\}$. There is a matrix R' of order two larger with the same determinant such that the selected row is replaced by binary entries, all added entries are binary, and every other old entry is unchanged. The nonunit Smith factors are also unchanged.*

Proof. Write $z = w - a - b$ with binary row vectors a, b, w . Coordinatewise, take

$$w_j = \mathbf{1}_{\{z_j=1\}}, \quad a_j = \mathbf{1}_{\{z_j<0\}}, \quad b_j = \mathbf{1}_{\{z_j=-2\}}.$$

Adjoin two first rows and two first columns. In the first two rows put the block $[I_2 \mid a; b]$. In the selected old row put $(1, 1)$ in the new columns and w in the old columns. In every other old row put zeros in the new columns and retain its old entries.

Elimination at I_2 replaces w by $w - a - b = z$ and leaves every other old row unchanged. These are integer unimodular operations, so R' is integrally equivalent to $I_2 \oplus R$. This proves both assertions and avoids any row-permutation sign convention. $\square$

Lemma III.78.4 (A bounded integer row). *Let C be an $n \times n$ integer matrix whose rows other than the first are binary. Suppose the entries of its first row have absolute value at most an integer $a \geq 1$, and put*

$$b = \lceil \log_2(a+1) \rceil.$$

There is a binary matrix C^ of order $n + 3b - 1$ with $\det C^* = \det C$. Its nonunit Smith factors are precisely those of C .*

Proof. Writing c_j for the first-row entries, choose signed binary digits $e_{ij} \in \{-1, 0, 1\}$ with

$$c_j = \sum_{i=1}^b 2^{b-i} e_{ij}.$$

They exist because $a \leq 2^b - 1$. Let $E = (e_{ij})$, and let C_0 consist of the bottom $n - 1$ rows of C . Define a $b \times (b - 1)$ matrix J by

$$J_{i,i} = 1, \quad J_{i+1,i} = -2 \quad (1 \leq i < b),$$

with all other entries zero. Form

$$C_1 = \begin{pmatrix} J & E \\ 0 & C_0 \end{pmatrix}.$$

Its order is $n + b - 1$. Add twice the first row to the second row and eliminate the first column at its pivot 1. Repeating this procedure combines the digit rows into the original first row of C ; explicitly, the new digit row after the first elimination is $2E_1 + E_2$. Integer column operations clear the rest of each pivot row. Thus C_1 is integrally equivalent to $I_{b-1} \oplus C$ and has determinant $\det C$.

Only the first b rows may contain nonbinary entries, all in $\{-2, -1, 0, 1\}$. Apply Lemma III.78.3 to each of those rows. This adds $2b$ to the order and adjoins only unit Smith factors. The resulting binary matrix has order $(n + b - 1) + 2b = n + 3b - 1$, as required. $\square$

The ceiling in this lemma is necessary, including when $a + 1$ is not a power of two. Its use is retained in every subsequent dimension count.

III.78.3 Substituting large entries at off-diagonal positions

Fix the data of Lemma III.78.2, an integer $t \geq 2$, and integers $\epsilon_0, \dots, \epsilon_{t-1}$. Let H be the $t \times t$ matrix whose first row is $(\epsilon_0, \dots, \epsilon_{t-1})$, with

$$H_{i,i} = 1, \quad H_{i,i-1} = \alpha \quad (2 \leq i \leq t),$$

and all remaining entries zero. Cofactor expansion in the first row gives

$$\det H = \sum_{j=0}^{t-1} (-\alpha)^j \epsilon_j. \quad (\text{III.78.4})$$

The following simultaneous substitution is useful because the large entries occur off the diagonal. It gives the determinant identity directly, without separate permutations of rows and columns.

Let H_0 be H with its $t - 1$ subdiagonal entries α replaced by zero, and let K be the block diagonal matrix with $t - 1$ copies of M . Define binary coupling matrices R, S as follows. For $i = 2, \dots, t$, row i of R contains $u^\top$ in auxiliary block $i - 1$ and is zero elsewhere; the first row of R is zero. Auxiliary block $i - 1$ of S contains v in column $i - 1$ and is zero elsewhere. Set

$$\tilde{H} = \begin{pmatrix} H_0 & R \\ S & K \end{pmatrix}.$$

Its order is

$$t + (t - 1)(q - 1) = (t - 1)q + 1. \quad (\text{III.78.5})$$

Because $\det K = 1$, elimination gives

$$\det \tilde{H} = \det(H_0 - RK^{-1}S) = \det H. \quad (\text{III.78.6})$$

Indeed, the only nonzero entries of $RK^{-1}S$ are $u^\top M^{-1}v = -\alpha$ in precisely the intended subdiagonal positions. Every row of $\tilde{H}$ other than the first is binary.

This substitution also preserves the integral normal form. The integer unimodular matrices

$$P = \begin{pmatrix} I & -RK^{-1} \\ 0 & I \end{pmatrix}, \quad Q = \begin{pmatrix} I & 0 \\ -K^{-1}S & I \end{pmatrix}$$

satisfy

$$P\tilde{H}Q = H \oplus K. \quad (\text{III.78.7})$$

Thus the added blocks contribute only unit Smith factors.

Proposition III.78.5 (An exact integer range at an exact order). *For every sufficiently large q , every integer $t \geq 2$, and every integer h with $|h| < A(q)^t$, there is a binary matrix of order*

$$(t-1)q + 3 \lceil \log_2 A(q) \rceil \quad (\text{III.78.8})$$

and determinant h . If $h \neq 0$, its Smith form has exactly one possibly nonunit factor, namely $|h|$.

Proof. Write $|h| = \sum_{j=0}^{t-1} A^j d_j$ in base A , with $0 \leq d_j \leq A-1$. For $h \neq 0$ set

$$\epsilon_j = \text{sgn}(h)(-\text{sgn}(\alpha))^j d_j;$$

for $h = 0$ take every $\epsilon_j = 0$. Equation (III.78.4) gives $\det H = h$. Substitute the auxiliary blocks to obtain $\tilde{H}$, and apply Lemma III.78.4 with $a = A-1$ and $b = \lceil \log_2 A \rceil$. Combining its cost with (III.78.5) gives exactly (III.78.8).

The bottom $t-1$ rows and right $t-1$ columns of H form a triangular matrix with diagonal 1. Hence the first $t-1$ Smith factors of H are 1, and its last factor is $|h|$ when $h \neq 0$. Equation (III.78.7) and Lemma III.78.4 show that every subsequent operation adjoins only unit factors. $\square$

Remark III.78.6 (A single unit cofactor). The stronger cofactor statement in Theorem III.78.1 can also be tracked directly. Deleting the first row and first column of $\tilde{H}$ gives determinant 1: its Schur complement is the unit triangular minor just used, and $\det K = 1$. In the digit-encoding matrix C_1 , delete its last digit row and the original first column. The remaining upper-left $(b-1) \times (b-1)$ block is triangular with diagonal 1, and the lower-right block is the previous unit minor. Each two-row conversion preserves that minor by retaining the new I_2 block and omitting the same old row and column. Identity padding preserves it as well. This gives an actual minor equal to 1, rather than only a greatest common divisor of minors equal to 1.

III.78.4 An improved secondary term

Proof of Theorem III.78.1. For the target order N , write $L = \log N$ and set

$$q = \left\lceil \frac{N}{L^2} \right\rceil, \quad t = \lfloor L^2 - 3L \rfloor, \quad b = \lceil \log_2 A(q) \rceil. \quad (\text{III.78.9})$$

For all sufficiently large N , the integer q lies in the range of Lemma III.78.2, $t \geq 2$, $N/L^2 > e^2$, and $q \leq N$. Put $c = 3/(2 \log 2) < 2.2$. The upper bound in (III.78.3) gives $3b \leq cq \log q + 3$. Consequently the order s in (III.78.8) satisfies

$$\begin{aligned} s &\leq q[L^2 - 3L - 1 + c \log q] + 3 \\ &\leq q[L^2 - (3 - c)L - 1] + 3 \\ &\leq N - \frac{(3 - c)N}{L} - \frac{N}{L^2} + L^2 - (3 - c)L + 2 \\ &\leq N - \frac{(3 - c)N}{L} + L^2 + 2 \leq N. \end{aligned}$$

The coefficient of q is positive eventually, justifying the insertion $q \leq N/L^2 + 1$. The last inequality holds eventually because $3 - c > 0$ and N/L grows faster than L^2 . A matrix of order s can therefore be padded by an identity matrix to exact order N .

For the covered range, the function $x \mapsto \frac{x}{2} \log x - \frac{3}{2}x$ is increasing for $x > e^2$. Since $q \geq N/L^2$, the lower bound in (III.78.3) gives

$$\log A(q) \geq \frac{N}{L^2} \left(\frac{L}{2} - \log L - \frac{3}{2} \right).$$

Also $t \geq L^2 - 3L - 1$. Both lower bounds are positive eventually, and their product is

$$\begin{aligned} t \log A(q) &\geq N \left[\frac{L}{2} - \log L - 3 + \frac{3 \log L + 4}{L} + \frac{\log L + 3/2}{L^2} \right] \\ &> \frac{N}{2} \log N - N \log \log N - 3N. \end{aligned}$$

In particular, the exponential in (III.78.1) is strictly less than $A(q)^t$. Every integer in the stated interval is covered by Proposition III.78.5. Its Smith-form assertion, Remark III.78.6, and identity padding prove all the remaining claims. $\square$

Remark III.78.7 (Quantifiers and attribution). The theorem is eventual: no numerical starting value N_0 is asserted. The threshold includes that in the asymptotic Alon–Vũ input. The auxiliary-block and radix strategy is credited to Shitov [147]; the displayed secondary estimate and the integral normal-form conclusions follow from the explicit calculations above. The factor 4 is a convenient stated constant, not an optimality claim.

Proposition III.78.8 (Optimizing the linear loss). *Put*

$$c = \frac{3}{2 \log 2}, \quad C_* = 2 \log 2 + \frac{1}{2} \log c + \frac{1}{2} = 2.272283375 \dots$$

For every $\varepsilon > 0$, Theorem III.78.1, including its Smith-form and unit-minor conclusions, holds with $4N$ in (III.78.1) replaced by $(C_ + \varepsilon)N$, for all sufficiently large N depending on ε .*

Proof. Retaining the sharper asymptotic input (III.78.2) in Lemma III.78.2 gives

$$\log A(q) \geq \frac{q}{2} \log q - (2 \log 2 + o(1))q.$$

Write $L = \log N$ and choose

$$q = \left\lceil \frac{N}{cL^2} \right\rceil, \quad t = \left\lfloor \frac{N - 3}{q} - c \log q \right\rfloor + 1.$$

Eventually $t \geq 2$. The unchanged estimate $3\lceil \log_2 A(q) \rceil \leq cq \log q + 3$ shows directly that the order in (III.78.8) is at most N . Moreover,

$$\log q = L - 2 \log L - \log c + o(1), \quad q(\log q)^2 = \frac{N}{c} + o(N).$$

Using $t \geq (N-3)/q - c \log q$ and the lower bound for $\log A(q)$, whose factors are positive eventually, yields

$$\begin{aligned} t \log A(q) &\geq \frac{N}{2} \log q - (2 \log 2)N - \frac{c}{2} q(\log q)^2 + o(N) \\ &= \frac{N}{2} \log N - N \log \log N - C_* N + o(N). \end{aligned}$$

Here $q \log q = o(N)$ absorbs the mixed term and the fixed rounding allowances. The positive margin εN absorbs the final $o(N)$ term and makes the proposed range strictly smaller than $A(q)^t$. Apply Proposition III.78.5 and pad by an identity block as before. $\square$

The choice of c minimizes this particular asymptotic parameter bound: using $q \sim N/(kL^2)$ instead gives linear loss $2 \log 2 + \frac{1}{2} \log k + c/(2k)$, minimized at $k = c$. This is not an optimality assertion for the determinant range itself.

III.79 Asymptotically optimal support for cyclic torsion

Source 34, delivered Section 7.

Recall the cyclic support parameter of Definition III.74.3. The support-compression lemma (Lemma III.75.5) transfers Theorem III.78.1 to a Boolean configuration, including composite cyclic orders. No assertion that all nonzero residues are units is needed.

Theorem III.79.1 (Leading asymptotic and an explicit eventual upper bound). *As $m \rightarrow \infty$ through the integers,*

$$h_{\text{cyc}}(m) \sim \frac{2 \log m}{\log \log m}. \quad (\text{III.79.1})$$

More precisely, put

$$L = \log m, \quad \lambda = \log L, \quad \mu = \log \lambda.$$

For every sufficiently large m ,

$$h_{\text{cyc}}(m) \leq \frac{2L}{\lambda - 3\mu + \log 2 - 8} + 2. \quad (\text{III.79.2})$$

For primes $p \rightarrow \infty$, the corresponding first-support parameter obeys

$$h(p) \sim h_{\text{cyc}}(p) \sim \frac{2 \log p}{\log \log p}, \quad D_{\text{div}}(p) \sim \frac{2 \log p}{\log \log p}.$$

Proof. Set $D = \lambda - 3\mu + \log 2 - 8$ and $n_0 = 2L/D$. For sufficiently large m , one has $2 \leq D \leq \lambda$, so

$$\log n_0 = \lambda + \log 2 - \log D \geq \lambda - \mu + \log 2, \quad \log \log n_0 \leq \mu.$$

It follows that

$$\frac{n_0}{2} (\log n_0 - 2 \log \log n_0 - 8) \geq \frac{n_0 D}{2} = L.$$

The function $x \mapsto \frac{x}{2} \log x - x \log \log x - 4x$ is increasing for all sufficiently large x . Thus $N = \lceil n_0 \rceil$ lies in the range of Theorem III.78.1 eventually, and that theorem supplies a binary matrix of order N with cyclic Smith factor m . Lemma III.75.5 gives $h_{\text{cyc}}(m) \leq N + 1 \leq n_0 + 2$, proving (III.79.2).

For the matching lower bound, Lemma III.76.1 gives

$$m \leq \frac{s^{s/2}}{2^{s-1}} \quad (\text{III.79.3})$$

for every admissible cyclic support of size s .

Taking logarithms in (III.79.3) gives $\log m \leq \frac{s}{2} \log s - (s-1) \log 2$. Together with the proved upper bound, this implies (III.79.1). For a prime p , one has $h(p) \leq h_{\text{cyc}}(p)$, and the determinant characterization $h(p) = D_{\text{div}}(p) + 1$ gives the same Hadamard lower bound. The remaining asymptotic statements follow. $\square$

Remark III.79.2 (The second-order gap remains). The theorem determines the leading constant, not a full asymptotic expansion of either minimum. For $m > 2$, the Hadamard comparison can be inverted exactly as

$$h_{\text{cyc}}(m) \geq \frac{2(L - \log 2)}{W((L - \log 2)/2)},$$

where W is the nonnegative real solution of $W(z)e^{W(z)} = z$. The lower comparison quantity expands as

$$\frac{2L}{\lambda} \left(1 + \frac{\mu + \log 2}{\lambda} + O\left(\frac{\mu^2}{\lambda^2}\right) \right),$$

whereas the upper comparison in (III.79.2) expands as

$$\frac{2L}{\lambda} \left(1 + \frac{3\mu - \log 2 + 8}{\lambda} + O\left(\frac{\mu^2}{\lambda^2}\right) \right).$$

The coefficients of μ/λ in these two comparison quantities are 1 and 3. Closing that gap is a further problem; neither coefficient is asserted to be the second-order coefficient of the actual minimum.

III.80 Universal positivity and its exact asymptotic cost

Source 34, delivered Section 8.

For $m \geq 2$, define $h_{\text{pos}}(m)$ to be the smallest number of vertices of an affinely independent Boolean configuration S such that $A_S \cong \mathbb{Z}/m\mathbb{Z}$, a faithful generator-coordinate vector $c = (c_1, \dots, c_s)$ has full support, and

$$T_S(f) \geq 0 \quad \text{for every finite abelian group } G \text{ and every real centered function } f : G \rightarrow \mathbb{R}.$$

No pairing of the coordinates is assumed in this definition. The following characterization shows that universal positivity forces precisely the pairing used in the constructive detector.

Theorem III.80.1 (Universal positivity is equivalent to inverse pairing). *Let S be an affinely independent Boolean configuration with $A_S \cong \mathbb{Z}/m\mathbb{Z}$ and full-support faithful generator coordinates $c_1, \dots, c_s \in \mathbb{Z}/m\mathbb{Z}$. The following are equivalent:*

1. $T_S(f) \geq 0$ for every finite abelian group and every real centered f .

2. The same statement holds just for $G = \mathbb{Z}/m\mathbb{Z}$.

3. The multiset $\{c_1, \dots, c_s\}$ can be partitioned into inverse pairs $\{a, -a\}$.

In the third condition, a nonzero self-inverse element, when present, must have even multiplicity. In particular, $s = 2t$.

Proof. The first condition implies the second. If the coordinates can be reordered as $(a_1, \dots, a_t, -a_1, \dots, -a_t)$, the Fourier formula of Proposition III.74.2 gives

$$T_S(f) = \sum_{\substack{\chi \in \widehat{G} \\ \chi^m = 1}} \prod_{j=1}^t |\widehat{f}(\chi^{a_j})|^2 \geq 0, \quad (\text{III.80.1})$$

proving the first condition.

For necessity, identify the dual of $\mathbb{Z}/m\mathbb{Z}$ with $\mathbb{Z}/m\mathbb{Z}$. Choose independent unit complex phases z_r for the non-self-inverse pairs, with $z_{-r} = \overline{z_r}$. If m is even, choose the additional phase $z_{m/2}$ independently from $\{1, -1\}$. The phase space is a compact torus, times a two-element group when necessary, with its probability Haar measure. Let $o(r)$ be the additive order of r . For a parameter $R > 0$, prescribe Fourier coefficients

$$\widehat{f}_R(0) = 0, \quad \widehat{f}_R(r) = R^{o(r)} z_r \quad (r \neq 0).$$

Fourier inversion gives a real centered function. Its moment is

$$T_S(f_R) = \sum_{u \in \mathbb{Z}/m\mathbb{Z}} \prod_{j=1}^s \widehat{f}_R(uc_j). \quad (\text{III.80.2})$$

Put $D = \sum_j o(c_j)$. If u is a unit modulo m , its summand has degree D in R . If u is not a unit, the summand is either zero or has strictly smaller degree. Indeed, multiplication cannot increase the order of any element. Choose a prime dividing $\gcd(u, m)$. Faithfulness means $\gcd(m, c_1, \dots, c_s) = 1$, so some c_j has full order at this prime; multiplication by u strictly decreases that order. Consequently,

$$\lim_{R \rightarrow \infty} R^{-D} T_S(f_R) = H(z) := \sum_{u \in (\mathbb{Z}/m\mathbb{Z})^\times} \prod_{j=1}^s z_{uc_j}. \quad (\text{III.80.3})$$

Suppose the coordinate multiset cannot be partitioned into inverse pairs. Then $\prod_j z_{c_j}$ is a nontrivial character of the phase group: either an inverse pair has unequal multiplicities, or the self-inverse element has odd multiplicity. Multiplication by a unit permutes inverse pairs and preserves this failure. Every summand of H is therefore a nontrivial character, and $\int H = 0$. On the other hand, $H(1) = \varphi(m) > 0$. The function H is real, by pairing u with $-u$; it is also continuous. These facts force $H(z) < 0$ for some phase choice. For that choice, (III.80.3) gives $T_S(f_R) < 0$ for all sufficiently large R , contradicting the second condition. This proves necessity for every m , including composite and even orders. $\square$

Lemma III.80.2 (Pairing any cyclic presentation). *If a full-support cyclic configuration has s vertices in dimension d , there is a universally positive configuration with the same cyclic quotient, $2s$ vertices, and dimension $d + s$. Consequently,*

$$h_{\text{cyc}}(m) \leq h_{\text{pos}}(m) \leq 2h_{\text{cyc}}(m). \quad (\text{III.80.4})$$

Proof. Let M be the original augmented Boolean matrix. Use generators $x_1, \dots, x_s, y_1, \dots, y_s$, retain the original rows on the x coordinates only, and add the binary relations $x_j + y_j = 0$. Eliminating $y_j = -x_j$ by integer unit pivots preserves the cyclic quotient and gives full-support coordinates $(c, -c)$. Replace the first pair row by the sum of all pair rows. This unimodular row operation produces an all-ones row, so the remaining $d + s$ binary rows define a Boolean configuration. Its finite quotient ensures rational column rank $2s$, hence distinctness and affine independence of the vertices. Formula (III.80.1) proves positivity. $\square$

Theorem III.80.3 (The sharp leading constant for positive detectors). *Every universally positive cyclic configuration of order m has an even number $s = 2t$ of vertices and satisfies*

$$m \leq t^{t/2}. \quad (\text{III.80.5})$$

As $m \rightarrow \infty$ through all positive integers,

$$h_{\text{pos}}(m) \sim \frac{4 \log m}{\log \log m}, \quad \frac{h_{\text{pos}}(m)}{h_{\text{cyc}}(m)} \rightarrow 2. \quad (\text{III.80.6})$$

Thus universal positivity has an asymptotic factor-two cost in the number of vertices. The explicit lower comparison is

$$h_{\text{pos}}(m) \geq \frac{4 \log m}{W(2 \log m)},$$

with W as in Remark III.79.2.

Proof. By Theorem III.80.1, order the generator coordinates as $(a, -a)$, with $a \in (\mathbb{Z}/m\mathbb{Z})^t$. Let $L = \text{row}_{\mathbb{Z}}(M) \subseteq \mathbb{Z}^{2t}$. Exactness of the quotient and faithfulness give

$$L = \ker((x, y) \mapsto a \cdot (x - y) \pmod{m}).$$

In particular, L contains the pair rows $e_j + e_{t+j}$. Project by $\pi(x, y) = x - y$. The quotient $\mathbb{Z}^t / \pi(L)$ is still $\mathbb{Z}/m\mathbb{Z}$, and the projections of the original binary rows are ternary rows, with entries in $\{-1, 0, 1\}$. They span a rank- t lattice. Choose t independent projected rows and place them in a square matrix B . Its nonzero determinant is a multiple of m , because its row lattice is a sublattice of $\pi(L)$, whose index is m . Every row has Euclidean norm at most $\sqrt{t}$. Hadamard's inequality proves $m \leq |\det B| \leq t^{t/2}$.

Inverting this inequality gives $t \geq (2 - o(1)) \log m / \log \log m$, hence the required lower bound with constant four. The cyclic construction of Theorem III.79.1, together with Lemma III.75.5, supplies $h_{\text{cyc}}(m) \sim 2 \log m / \log \log m$. Lemma III.80.2 gives the matching upper bound and then the ratio limit in (III.80.6). $\square$

Corollary III.80.4 (The first positive support sizes). *One has*

m	2	3	4	5	6	7	8
$h_{\text{pos}}(m)$	4	6	6	8	8	8	8.

Proof. The largest absolute determinants of ternary matrices of orders 1, 2, 3 are respectively 1, 2, 4. For order three, the maximum of the absolute determinant over the cube $[-1, 1]^9$ is attained at a sign matrix, by optimizing its entries one at a time. Such a determinant is divisible by four, while Hadamard bounds its absolute value by $3\sqrt{3} < 8$; hence the maximum is four. The smaller orders are immediate. Applying the ternary reduction in the proof of Theorem III.80.3 gives the displayed lower bounds. The elementary construction in Theorem III.77.1 has $2\lceil \log_2 m \rceil + 2$ vertices, which attains each bound. $\square$

III.81 The exact quartic term and its interface with Gowers's lemma

Source 34, delivered Section 9.

To connect the algebraic configurations and the norm comparison to the original paper, we recall the exact quartic expansion and give a short proof. This formula was already proved in the repository [155]; its role here is to make the subsequent improvement of the coefficient self-contained.

For a real function $F : G \rightarrow \mathbb{R}$, write

$$Q_d(F) = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{v \in \{0,1\}^d} F(x + v \cdot h).$$

This is $\|F\|_{U^d}^{2d}$. For an indicator function on a group of order N , the unnormalized number of parameterized cubes is $N^{d+1}Q_d(F)$. Degenerate choices of steps are included, as in Gowers's cube count.

Lemma III.81.1 (Small Boolean supports on odd-order groups). *Let G have odd order and let $f : G \rightarrow \mathbb{R}$ be centered. For a Boolean support S , the moment $T_S(f)$ is zero when $1 \leq |S| \leq 3$. At $|S| = 4$, it is zero unless S is a parallelogram of four distinct vertices over the integers. In the parallelogram case it equals $S_4(f) := \sum_{\chi} |\hat{f}(\chi)|^4$.*

Proof. Every two distinct Boolean vertices have a coordinate separating them, so their augmented matrix has a unit minor. Three distinct Boolean vertices are affinely independent: a straight line can meet the vertices of a cube in at most two points. A full minor including the all-ones row reduces to a nonsingular binary 2×2 determinant, necessarily of absolute value one. Thus the row quotient is trivial for supports of size at most three. Proposition III.74.2 and centering give vanishing.

For four affinely independent vertices, a nonsingular full minor reduces to a binary 3×3 determinant of absolute value at most two, by Lemma III.76.1. The row quotient has order one or two. It admits only the trivial homomorphism into the character group of an odd-order group, so the moment again vanishes.

If four distinct Boolean vertices are affinely dependent, their augmented rank is three. Any three of them give a unit minor, so the row quotient is torsion-free of rank one. Their primitive affine relation has two positive and two negative coefficients: one coefficient on either side would express a cube vertex as a strict convex combination of other cube vertices, which is impossible. Write the positive coefficients as a, b and the negative ones as $-c, -e$, so $a + b = c + e$.

In a coordinate where the four bits are not all equal, neither side can use both positive coefficients against only one negative coefficient: the remaining coefficient would then be zero. Hence the nonconstant coordinate pairs one coefficient on each side, with equal values. If the two positive coefficients were unequal, all such coordinate matchings would be forced, making a vertex on one side identical to a vertex on the other side. Therefore all four coefficients are equal, and primitivity gives the relation $(1, 1, -1, -1)$ after ordering. This is exactly an integer parallelogram. The quotient is $\mathbb{Z}$ with these generator coordinates. Its Fourier moment is $\sum_{\chi} \hat{f}(\chi)^2 \hat{f}(\chi^{-1})^2 = S_4(f)$. $\square$

Proposition III.81.2 (Exact quartic expansion, recalled with proof). *Let $d \geq 2$, let G have odd order, and let $F = \delta + f$ be real with $\delta \geq 0$ and $\mathbb{E}f = 0$. Put $u = \|f\|_{U^d}$ and*

$$P_d = \frac{6^d - 2 \cdot 4^d + 2^d}{8}.$$

Then

$$Q_d(\delta + f) = \delta^{2^d} + P_d \delta^{2^d-4} \|f\|_{U^2}^4 + R_d, \quad (\text{III.81.1})$$

where

$$|R_d| \leq \sum_{j=5}^{2^d} \binom{2^d}{j} \delta^{2^d-j} u^j. \quad (\text{III.81.2})$$

An empty sum is zero, and a zeroth power of δ is one.

Proof. Expand the product by its selected support. Lemma III.81.1 handles all supports of size at most four. To count the four-vertex parallelograms, first count ordered solutions of $v_1 + v_2 = v_3 + v_4$ in the Boolean cube over the integers. Each coordinate has six admissible bit patterns, giving 6^d solutions. A repeated vertex forces one of the two pairwise identifications $(v_1, v_2) = (v_3, v_4)$ or $(v_1, v_2) = (v_4, v_3)$. Each contributes 4^d , and their intersection contributes 2^d . The number with four distinct vertices is consequently $6^d - 2 \cdot 4^d + 2^d$. Each unordered parallelogram occurs in eight orderings, proving the formula for P_d .

For a support S of size $j \geq 5$, place f at its vertices and the constant one at all other vertices. Multilinear Gowers–Cauchy–Schwarz gives $|T_S(f)| \leq u^j$. Sum over these supports with their factors δ^{2^d-j} . $\square$

At $d = 4$, $P_4 = 100$. The sharp norm question is therefore exactly the question of the best universal leading constant in this perturbative cube bound when the error is measured by U^4 . The next section improves that constant. The earlier structural theorems answer a different question: at what support sizes and cube dimensions new prime Smith torsion can enter higher terms. For example, order seven can first occur at support degree seven in dimension six, while orders eleven and thirteen require support degree eight and dimension seven.

Remark III.81.3 (Compatibility with the original normalization). If $A \subseteq \mathbb{Z}/N\mathbb{Z}$ has density δ and $\|\mathbf{1}_A - \delta\|_{U^d}^{2^d} \leq \alpha$, insert $u \leq \alpha^{1/2^d}$ in (III.81.1)–(III.81.2) and multiply by N^{d+1} . This uses the same uniformity normalization as Gowers’s Lemma 3.10 after converting its sums to averages. The odd-order hypothesis is essential for the quartic formula used here; the earlier repository expansion has additional order-two terms on even-order groups.

III.82 Source 38’s Part I: exact quotient reduction in exponent five [write]

Added in the seventh write, 6 October 2026.

The next four sections are source 38’s Sections 2–5, its delivered Part I. *Logically they continue source 21’s chapter* (Section III.49) and source 23’s exponent-five theorem; they are printed here because no number printed earlier may move. Source 38’s introduction, with its main theorems, is Section F.43.

What is proved. On every nontrivial group of exponent five the maximum of the three-cube count of $\delta + f$ at fixed centered $\|f\|_{U^3} = u$ equals its value on $\mathbb{Z}/5\mathbb{Z}$ whenever $u \leq 10^{-5}\delta$, and every maximizer factors through $\mathbb{Z}/5\mathbb{Z}$ (Theorem III.84.1, with quantitative off-quotient control, Corollary III.84.2); so source 21’s u^8 coefficient 20/9 is universal, and the profile continues with u^{10} and u^{12} coefficients $869/(216 \cdot 2^{3/4})$ and $65759/(5832 \cdot 2^{3/2})$ (Theorem III.85.1); for every fixed odd exponent m the reduction is to a cyclic quotient when $u/\delta \leq 10^{-8}m^{-8}$ (Theorem III.86.1). This

answers three of source 21's questions (item III.27); source 38 credits sources 21 and 23 for $2^{3/4}/3$ and 20/9.

Relation to the formal statements (checked at HEAD 89c1cbc38). `lemma_3_10` is exact (companion `lemma_3_10_holds`, `Proofs03CubeUpper.lean`, line 82); its scope is all moduli, and a theorem on $\mathbb{F}_5^n$ does not replace it, as source 38 says. Nothing in this chapter is formalized.

III.83 Fourier identities and the retained energy defect

Source 38, delivered Section 2.

*Delivered part heading of source 38: **Exact characteristic-five cube extremizers***

This section records the background identities in the form needed by the new projection argument. The fourth-order expansion and its leading norm gap were established in preceding project reports [169, 170]; they are reproved here so the main result has no dependence on an unquoted computation.

For $f : G \rightarrow \mathbb{C}$, use

$$a_r = \widehat{f}(r) = \mathbb{E}_x f(x) \overline{r(x)}, \quad f(x) = \sum_{r \in \widehat{G}} a_r r(x).$$

For real centered f , one has $a_{-r} = \overline{a_r}$ and $a_0 = 0$. Set

$$\lambda_r = |a_r|^2, \quad S = \sum_r |a_r|^4, \quad M = \sum_r |a_r|^8, \quad \mathcal{E}(\lambda) = \sum_{r+s=t+w} \lambda_r \lambda_s \lambda_t \lambda_w.$$

The Gowers–Cauchy–Schwarz inequality bounds a cube with eight possibly different functions by the product of their U^3 norms. It follows by applying Cauchy–Schwarz successively in the three directions and expanding each resulting square [1, Lemma 3.8]. Its homogeneous form requires no L^∞ hypothesis. Expanding a sum in the cube form then gives the triangle inequality. In particular, translation averages are contractive for U^3 .

Lemma III.83.1 (Fourier squares and their axes). *Put*

$$T_f(s, t) = \sum_r a_r \overline{a_{r+s}} \overline{a_{r+t}} a_{r+s+t}.$$

Then

$$\|f\|_{U^3}^8 = \sum_{s,t} |T_f(s, t)|^2.$$

For real f , T_f is real, symmetric in its two arguments, and unchanged by reversing either argument. Its two coordinate axes, with their common origin counted once, contribute exactly

$$2\mathcal{E}(\lambda) - S^2.$$

Proof. Expand $|T_f(s, t)|^2$ and sum over s, t . Fourier expansion of the physical cube imposes the base and three direction equations on its eight frequencies. Their solutions are parametrized by a parity-adjusted affine frequency cube; grouping its two opposite faces gives precisely the displayed sum of squares. This can also be checked directly by using the two base frequencies and the two increments in the expanded square as its four free variables. No division by a group element is used.

For real f , substitution $r \mapsto -r - s - t$ pairs every summand in T_f with its conjugate. Interchange of s, t and translation of r give the other symmetries. Finally

$$T_f(0, t) = \sum_r \lambda_r \lambda_{r+t}, \quad T_f(0, 0) = S.$$

The squared sum on either axis is $\mathcal{E}(\lambda)$; subtract one copy of S^2 at their intersection. $\square$

Lemma III.83.2 (Odd-order quartic gap). *For real centered f on a finite abelian group of odd order,*

$$\|f\|_{U^3}^8 \geq 5S^2 - 6M \geq 2S^2.$$

If $\{\xi, -\xi\}$ is a largest conjugate Fourier pair and $D_{\text{pair}} = \sum_{r \notin \{\xi, -\xi\}} |a_r|^4$, then

$$\|f\|_{U^3}^8 \geq 2S^2 + 3SD_{\text{pair}}.$$

A pure conjugate pair whose character has order five attains $\|f\|_{U^3}^8 = 2S^2$.

Proof. Write the energy as a sum over zero-sum quadruples and retain the three families obtained by pairing the four positions into opposite pairs. Each family has weight S^2 , and every pairwise intersection has weight M . The triple intersection requires $2r = 0$; it has zero weight because the order is odd and $a_0 = 0$. Thus $\mathcal{E}(\lambda) \geq 3S^2 - 3M$. Retaining the axes in Lemma III.83.1 proves the first inequality.

If x_j are the total fourth-power masses of the conjugate pairs, then $S = \sum_j x_j$ and $M = \frac{1}{2} \sum_j x_j^2 \leq S^2/2$. For a largest pair with coefficient modulus ρ , $M \leq \rho^4 S = (S - D_{\text{pair}})S/2$, proving the stronger bound. For a pure pair of order five, the only nonzero Fourier-square slots are $(0, 0)$ and the four axis slots with increment $\pm 2\xi$. They contribute $4\rho^8 + 4\rho^8 = 8\rho^8 = 2S^2$. $\square$

Lemma III.83.3 (Complete centered expansion through degree seven). *Assume G has odd order and f is real and centered. Then*

$$\mathcal{C}_3(\delta + f) = \delta^8 + 12\delta^4 S + 8\delta^3 B_5(f) + \delta^2 C_6(f) + \delta C_7(f) + \|f\|_{U^3}^8,$$

where

$$B_5(f) = \sum_r a_{2r} a_{-r}^3 a_r, \quad C_6(f) = 12H_f + 12J_f + 4K_f,$$

with

$$\begin{aligned} H_f &= \sum_{r+s+t=0} \lambda_r \lambda_s \lambda_t, \\ J_f &= \sum_{r,s} \lambda_r a_s^2 a_{r-s} \overline{a_{r+s}}, \\ K_f &= \sum_{r+s+t=0} a_r^2 a_s^2 a_t^2. \end{aligned}$$

These statistics are real. If

$$L_f(s, t) = 2 \operatorname{Re}(a_s a_t \overline{a_{s+t}}) + 2 \operatorname{Re}(a_s \overline{a_t} \overline{a_{s-t}}),$$

then $C_7(f) = 2 \operatorname{Re}\langle T_f, L_f \rangle$, with counting measure in this inner product. In particular

$$|B_5(f)| \leq \|f\|_{U^3}^5, \quad |C_6(f)| \leq 28\|f\|_{U^3}^6, \quad |C_7(f)| \leq 8\|f\|_{U^3}^7.$$

Proof. Expand the eight-vertex product by the support of its centered factors. A support of size one, two, or three consists of independent cube forms whose coefficient equations have unit pivots, so centering kills it. For four vertices, the twelve parallelogram supports contribute S each; all other supports are independent. For five vertices, forty-eight supports force a zero Fourier coefficient. The other eight are cube symmetries of $\{000, 100, 010, 001, 111\}$, whose relation has coefficients $(2, -1, -1, -1, 1)$; their contribution is B_5 .

For clarity, these assertions follow by solving the four equations

$$\sum_{\omega} r_{\omega} = 0, \quad \sum_{\omega} \omega_j r_{\omega} = 0 \quad (1 \leq j \leq 3),$$

with frequencies outside the chosen support set equal to zero. All the relation systems here can be solved using only unit pivots and division by two. The latter is available in every group of odd order; in particular, these identities do not require exponent five. For six vertices classify the two omitted vertices by their Hamming distance:

omitted representative	number	contribution
000, 001	12	H_f
000, 011	12	J_f
000, 111	4	K_f .

Solving the same equations gives the displayed formulas. Their reality also follows from the real spatial cube products.

Replace a_0 by a real variable and differentiate the Fourier-square identity at zero. Differentiating its four factors gives exactly the four cubic summands in L_f ; differentiating the squared norm gives $2 \operatorname{Re}\langle T_f, L_f \rangle$. Repeated indices cause no extra terms because the remaining zero coefficient then vanishes. Finally each support is bounded by cube Cauchy–Schwarz, with constants at its complementary vertices. There are 28 six-vertex and 8 seven-vertex supports. $\square$

The new argument will use these exact formulas to compare a function with its projection. Bounding C_7 only by the product of norms would lose a square root of the projection deficit. The exact derivative formula avoids that loss.

III.84 Exact reduction to a five-point quotient

Source 38, delivered Section 3.

Write $N(f) = \|f\|_{U^3}^8 = C_3(f)$ for real functions. All Fourier transforms are normalized, while frequency sums are unnormalized. The following proves more than the formerly open universal eighth-order coefficient: at sufficiently small centered norm, the entire extremal problem reduces exactly to $\mathbb{Z}/5\mathbb{Z}$.

Theorem III.84.1 (Exact dimension reduction; explicit sufficient range). *Let $G \neq 0$ be a finite abelian group with $5G = 0$. For $\delta > 0$ and $0 < u \leq 10^{-5}\delta$, the maximum of $C_3(\delta + f)$ over real centered functions $f : G \rightarrow \mathbb{R}$ with $\|f\|_{U^3} = u$ is identical to the corresponding maximum on $\mathbb{Z}/5\mathbb{Z}$. Every maximizing f factors through a nonzero homomorphism $G \rightarrow \mathbb{Z}/5\mathbb{Z}$. The same conclusion for the maximum holds with norm at most u ; the unrestricted profile is strictly increasing throughout this range. For $0 < \delta < 1$, weights satisfying $0 \leq \delta + f \leq 1$ have the same maximum if, additionally, $4u \leq \min\{\delta, 1 - \delta\}$.*

[write] **6 October 2026; sources 34–39.** *Answers two of source 21's questions* (item III.27): the universal coefficient of u^8 exists in the stated range (for $u/\delta \leq 10^{-5}$ and interior weights) and is $20/9$, the value source 21 proves on $\mathbb{Z}/5\mathbb{Z}$ (its Corollary III.55.2), and it is the same on every fixed $\mathbb{F}_5^n$, with every maximizer factoring through $\mathbb{Z}/5\mathbb{Z}$. The cutoff 10^{-5} is a sufficient certificate; the reduction itself is not numerically testable at the scales the intake could afford (a 24-parameter probe on $(\mathbb{Z}/5\mathbb{Z})^2$ was stopped unfinished; no result is claimed from it).

The threshold 10^{-5} is only a convenient certificate and is not asserted sharp. We prove the norm-equality, unrestricted-real assertion first; norm-ball monotonicity and bounded Fourier support then give the last two statements.

Proof. Normalize $f = uF$, $\|F\|_{U^3} = 1$, and $t = u/\delta$. Thus the objective, after dividing by δ^8 , is $C_3(1 + tF)$. Let

$$a_r = \widehat{F}(r), \quad \lambda_r = |a_r|^2, \quad S = \sum_r |a_r|^4, \quad S_0 = 2^{-1/2}.$$

The exact centered expansion is

$$C_3(1 + tF) = 1 + 12St^4 + 8B_5(F)t^5 + C_6(F)t^6 + C_7(F)t^7 + t^8, \quad B_5(F) = \sum_r a_{2r}a_{-r}^3a_r.$$

The standard cube Cauchy–Schwarz inequalities give $|B_5(F)| \leq 1$, $|C_6(F)| \leq 28$, and $|C_7(F)| \leq 8$. The real centered norm gap gives

$$1 \geq 5S^2 - 6 \sum_r |a_r|^8 \geq 2S^2.$$

III.84.1 Localization requiring only the leading extremal term

A normalized pure Fourier pair gives exactly $1 + 6\sqrt{2}t^4 + t^8$. We may therefore assume the candidate F has at least this value. Set $\eta = S_0 - S$. For $t \leq 1$, the displayed expansion implies

$$0 \leq \eta \leq \frac{2}{3}t + \frac{7}{3}t^2 + \frac{2}{3}t^3 \leq 4t.$$

Choose a largest conjugate pair $\{\xi, -\xi\}$, put $a = |a_\xi|$, and let $D_{\text{pair}} = \sum_{r \notin \{\xi, -\xi\}} |a_r|^4$. Since $\sum_r |a_r|^8 \leq a^4 S = (S - D_{\text{pair}})S/2$,

$$1 \geq 2S^2 + 3SD_{\text{pair}}.$$

When $S \geq S_0/2$, this yields $D_{\text{pair}} \leq 2\eta \leq 8t$. Let $L = \langle \xi \rangle$ and let g be the Fourier projection of F to L , equivalently the subgroup average over $\ker \xi$. Write $h = F - g$, $b = |a_{2\xi}|$, and

$$S_g = \sum_{r \in L} |a_r|^4, \quad D = \sum_{r \notin L} |a_r|^4, \quad v = \|g\|_{U^3}, \quad \Delta = 1 - v^8.$$

Then

$$S_g \geq S_0 - 12t, \quad b \leq (4t)^{1/4}, \quad a \leq 2^{-3/8}.$$

Subgroup averaging is contractive for the Gowers norm, so $\Delta \geq 0$.

III.84.2 Coercivity of the norm loss outside the principal line

Set

$$T_F(s, t) = \sum_r a_r \overline{a_{r+s} a_{r+t}} a_{r+s+t}, \quad N(F) = \sum_{s, t} |T_F(s, t)|^2.$$

For $s, t \in L$, the base frequency r remains in a single coset of L throughout the parallelogram, and hence

$$T_F(s, t) = T_g(s, t) + T_h(s, t), \quad |T_h(s, t)| \leq D.$$

Let

$$R = \sum_{\substack{s, t \in L \\ s \neq 0, t \neq 0}} |T_g(s, t)|.$$

The elementary five-point slot enumeration gives

$$R \leq 8a^3b + 8a^2b^2 + 8ab^3.$$

Indeed the nonaxis slots, up to independent sign changes and interchange, are (ξ, ξ) with multiplicity four and modulus at most $2ab^3$, $(2\xi, 2\xi)$ with multiplicity four and modulus at most $2a^3b$, and $(\xi, 2\xi)$ with multiplicity eight and modulus a^2b^2 .

Let $\mathcal{E}(\lambda) = \sum_{r+s=t+w} \lambda_r \lambda_s \lambda_t \lambda_w$. The axes in the T sum contribute $2\mathcal{E}(\lambda) - S^2$. Write $\lambda_g = \lambda \mathbf{1}_L$ and $\lambda_h = \lambda \mathbf{1}_{L^c}$. In addition to all entirely inside and entirely outside energy tuples, retain the mixed tuples in the three families

$$r = t, s = w; \quad r = w, s = t; \quad s = -r, w = -t.$$

Each family has mixed weight $2S_g D$. Their pairwise intersections contain no mixed tuple, because L and its complement are stable under negation. All terms are nonnegative. Consequently

$$\mathcal{E}(\lambda) \geq \mathcal{E}(\lambda_g) + \mathcal{E}(\lambda_h) + 6S_g D.$$

The usual three-family argument on the even sequence λ_h gives $\mathcal{E}(\lambda_h) \geq 3D^2 - 3 \sum \lambda_h^4 \geq \frac{3}{2}D^2$; the last inequality uses conjugate pairing and odd group order. Subtracting $S^2 = (S_g + D)^2$ gives the axes estimate

$$2\mathcal{E}(\lambda) - S^2 \geq 2\mathcal{E}(\lambda_g) - S_g^2 + 10S_g D + 2D^2.$$

The nonaxis slots in L^2 lose at most $2RD$ relative to those of g , by $|T_g + T_h|^2 \geq |T_g|^2 - 2|T_g|D$; all other slots are nonnegative. We have therefore proved the exact finite inequality

$$\boxed{\Delta \geq (10S_g - 2R)D + 2D^2.} \quad (\text{III.84.1})$$

For $t \leq 10^{-5}$, we have $S_g \geq 2/3$ and the rational bounds $a < 193/250$ and $b < 2/25$. Indeed $8(193/250)^8 > 1$ and $(2/25)^4 > 4 \cdot 10^{-5}$. Direct rational substitution gives

$$8(193/250)^3(2/25) + 8(193/250)^2(2/25)^2 + 8(193/250)(2/25)^3 < 1/3.$$

Thus $R < 1/3$, $10S_g - 2R \geq 6$, and

$$D \leq \Delta/6, \quad v^8 \geq 2S_g^2 \geq 8/9, \quad 0 \leq \Delta \leq 1/9.$$

III.84.3 Quantitative subgroup projection

We use the general estimate

$$\|h\|_{U^3}^8 \leq 2^8 \Delta. \quad (\text{III.84.2})$$

For completeness, Gowers–Cauchy–Schwarz and the expansion of $\|f + g\|_{U^3}^8 + \|f - g\|_{U^3}^8$ give Clarkson’s inequality

$$\left\| \frac{f+g}{2} \right\|_{U^3}^8 + \left\| \frac{f-g}{2} \right\|_{U^3}^8 \leq \frac{\|f\|_{U^3}^8 + \|g\|_{U^3}^8}{2}.$$

To see the constant directly, the odd mixed terms cancel, each even mixed term is bounded by the product of the corresponding norms, and weighted arithmetic–geometric mean bounds the resulting scalar sum by $2^7(\|f\|_{U^3}^8 + \|g\|_{U^3}^8)$. Dividing by 2^8 gives the displayed inequality, in agreement with the general norm-geometric perspective of [31]. Apply it to F and each translate $T_y F$, $y \in \ker \xi$. Projection of their midpoint is g , so contractivity bounds the first term below by v^8 . This gives $\|F - T_y F\|_{U^3} \leq 2\Delta^{1/8}$; average over y and use the triangle inequality to obtain (III.84.2).

III.84.4 Sixth-degree differences are linear in the norm loss

Put

$$H_F = \sum_{r+s+t=0} \lambda_r \lambda_s \lambda_t.$$

The exact sixth-degree decomposition is

$$C_6(F) = 12H_F + 12J_F + 4K_F, \quad J_F = \sum_{r,s} \lambda_r a_s^2 a_{r-s} \overline{a_{r+s}}, \quad K_F = \sum_{r+s+t=0} a_r^2 a_s^2 a_t^2.$$

In the J sum, the entirely inside part is exactly the restriction $(r, s) \in L^2$. Cauchy–Schwarz on the complementary domain gives

$$|J_F - J_g| \leq H_F - H_g.$$

Indeed its two squared bounds are the sums of $\lambda_r \lambda_s \lambda_{r-s}$ and $\lambda_r \lambda_s \lambda_{r+s}$ over $(r, s) \notin L^2$; both equal $H_F - H_g$, using evenness and $s \mapsto -s$. Also $|K_F - K_g| \leq H_F - H_g$. Hence

$$|C_6(F) - C_6(g)| \leq 28(H_F - H_g).$$

A zero-sum triple cannot have exactly one frequency outside L . Writing $A_g = \sum_{r \in L} |a_r|^2$, note that g has at most four nonzero Fourier coefficients, so $A_g^2 \leq 4S_g \leq 4$ and $A_g \leq 2$. We obtain

$$H_F - H_g \leq 3A_g D + H_h.$$

The first term comes from the three choices of an inside frequency; for each such frequency, the remaining convolution is at most D by Cauchy–Schwarz. For the entirely outside term,

$$H_h \leq \sqrt{D\mathcal{E}(\lambda_h)} \leq \sqrt{D(128\Delta + D^2/2)},$$

because the axes inequality gives $\mathcal{E}(\lambda_h) \leq (\|h\|_{U^3}^8 + D^2)/2$ and (III.84.2) applies. Using $D \leq \Delta/6$ and $\Delta \leq 1$ therefore yields

$$H_F - H_g \leq 6\Delta, \quad |C_6(F) - C_6(g)| \leq 168\Delta. \quad (\text{III.84.3})$$

III.84.5 Seventh-degree differences are also linear in the norm loss

Define

$$L_F(s, t) = 2 \operatorname{Re}(a_s a_t \overline{a_{s+t}}) + 2 \operatorname{Re}(a_s \overline{a_t} \overline{a_{s-t}}).$$

Differentiating the sum-of-squares identity with respect to the zero Fourier coefficient gives $C_7(F) = 2 \operatorname{Re}\langle T_F, L_F \rangle$. For $s, t \in L$ one has $L_F(s, t) = L_g(s, t)$ exactly, because all its indices then belong to L . Thus the contribution to the difference from L^2 has modulus at most

$$2D \sum_{L^2} |L_g| \leq 200D \leq \frac{100}{3} \Delta.$$

Here $|a_r| \leq 1$, so each of the 25 values of L_g has modulus at most four.

For the complement of L^2 , the same expansion of the norm gives

$$\sum_{(s,t) \notin L^2} |T_F(s, t)|^2 \leq \Delta + 2RD \leq \frac{4}{3} \Delta.$$

The axes cross terms within L^2 are nonnegative, and the remaining cross terms lose at most $2RD$, which proves this inequality. Each of the four cubic summands of L_F has squared norm on the complement of L^2 equal to $H_F - H_g$. Consequently

$$\sum_{(s,t) \notin L^2} |L_F(s, t)|^2 \leq 16(H_F - H_g) \leq 96\Delta.$$

Cauchy–Schwarz bounds their contribution by $16\sqrt{2} \Delta$. Adding the two regions proves

$$|C_7(F) - C_7(g)| \leq 60\Delta. \quad (\text{III.84.4})$$

The linear estimate (III.84.4) is essential for exact localization; a generic mixed cube bound would give only $O(\sqrt{\Delta})$ and would settle the eighth-order coefficient without proving exact reduction.

III.84.6 Rescaling the projection strictly improves every off-line candidate

Doubling preserves L and its complement, so $B_5(F) - B_5(g) = B_5(h)$ exactly. Therefore

$$|B_5(F) - B_5(g)| \leq D^{5/4} \leq D \leq \Delta/6.$$

Let $g_* = g/v$; then $\|g_*\|_{U^3} = 1$. Since $0 \leq \Delta \leq 1/9 < 1/2$, for $k = 5, 6, 7$ the mean value theorem gives

$$|v^{-k} - 1| = |(1 - \Delta)^{-k/8} - 1| \leq 4\Delta.$$

Combining this with (III.84.3), (III.84.4), homogeneity, and cube Cauchy–Schwarz on g , one obtains

$$|B_5(F) - B_5(g_*)| \leq 5\Delta, \quad |C_6(F) - C_6(g_*)| \leq 280\Delta, \quad |C_7(F) - C_7(g_*)| \leq 92\Delta.$$

At fourth degree,

$$\begin{aligned} 12t^4(S - S_{g_*}) &= 12t^4[D - S_g((1 - \Delta)^{-1/2} - 1)] \\ &\leq 12t^4(D - S_g\Delta/2) \leq -2t^4\Delta. \end{aligned}$$

The eighth-degree terms cancel exactly because F and g_* have the same norm. Hence

$$\mathcal{C}_3(1 + tF) - \mathcal{C}_3(1 + tg_*) \leq t^4\Delta(-2 + 40t + 280t^2 + 92t^3) \leq -t^4\Delta \quad (0 < t \leq 10^{-5}). \quad (\text{III.84.5})$$

If F has any nonzero off-line coefficient, then $D > 0$, and (III.84.1) implies $\Delta > 0$. Thus its projected and rescaled competitor is strictly better. If the candidate was initially below the pure-pair value, the pure-pair competitor already improves it. Every maximum therefore occurs on a five-point quotient. Conversely, every function on the five-point group pulls back along any nontrivial character and preserves its mean, all Gowers norms, and its cube count. This proves the equality of maxima and the assertion about every maximizer.

For completeness, the norm-ball conclusion does not require an unspecified smaller analytic neighborhood. Let F be a normalized maximizer at norm $0 < r \leq 10^{-5}$. Localization gives $S \geq 2/3$. For $r \leq s \leq 10^{-5}$, direct differentiation of the centered polynomial gives

$$\frac{d}{ds}\mathcal{C}_3(1 + sF) \geq s^3(48S - 40s - 168s^2 - 56s^3) > 0.$$

Thus scaling a lower-norm maximizer strictly increases its objective throughout the range. The norm-zero value is strictly smaller by the pure-pair benchmark. The norm-equality maximum is strictly increasing, and the norm-ball maximum occurs on its boundary.

For interval-valued weights, every Fourier coefficient of a maximizer has modulus at most its U^3 norm u . Its support has at most four frequencies, so $\|f\|_\infty \leq 4u$. Thus the unrestricted maximizing function is feasible for the interval constraint in the stated range; the reverse inequality between maxima is immediate from inclusion. $\square$

III.84.7 Quantitative stability of quotient dependence

Corollary III.84.2 (Off-quotient control by the extremal deficit). *Assume the hypotheses of Theorem III.84.1 and that f has at least the pure-pair benchmark,*

$$\mathcal{C}_3(\delta + f) \geq \delta^8 + 6\sqrt{2}\delta^4u^4 + u^8.$$

Let $\mathcal{D} = \mathcal{M}_{\delta,G}(u) - \mathcal{C}_3(\delta + f)$. There is a nonzero character ξ , with P the average over its kernel, such that

$$\|f - Pf\|_{U^2}^4 \leq \frac{\mathcal{D}}{6\delta^4}, \quad (\text{III.84.6})$$

$$\|f - Pf\|_{U^3}^8 \leq \frac{256u^4\mathcal{D}}{\delta^4}. \quad (\text{III.84.7})$$

The constants do not depend on G .

Proof. Apply the construction in the proof with $F = f/u$. The rescaled projection is an admissible competitor, so (III.84.5) gives $\mathcal{D} \geq \delta^8 t^4 \Delta = \delta^4 u^4 \Delta$. Now use $D \leq \Delta/6$, the Fourier identity $\|f - Pf\|_{U^2}^4 = u^4 D$, and (III.84.2). $\square$

For fixed $\delta > 0$, attaining the extremal expansion with error $o(u^8)$ therefore forces an off-quotient U^2 error $o(u^2/\delta)$ and an off-quotient U^3 error $o(u^{3/2}/\delta^{1/2})$. The benchmark assumption is automatic for such a family at sufficiently small u , because the sharp u^6 coefficient is positive. The conclusion concerns uniformity norms. It does not assert that indicator approximations have vanishing L^2 distance from a nonconstant fractional weight.

III.85 The five-point profile and its analytic expansion

Source 38, delivered Section 4.

Theorem III.84.1 reduces the dimension-uniform extremal problem to a fixed, four-dimensional real space of centered functions. We now solve its local optimization, including the phase variables. The coefficient 20/9 was calculated for this fixed group in the preceding characteristic-five report [170]; it is reproduced below as part of the calculation of two additional terms.

Write $\mathcal{M}_{\delta,5}(u) = \mathcal{M}_{\delta,\mathbb{Z}/5\mathbb{Z}}(u)$. The interval restriction $0 \leq \delta + f \leq 1$ is temporarily absent. This permits the exact scaling relation

$$\mathcal{M}_{\delta,5}(u) = \delta^8 \mathcal{M}_{1,5}(u/\delta). \quad (\text{III.85.1})$$

Theorem III.85.1 (The extremal profile through twelfth order). *There are absolute constants $\eta, C > 0$ such that, for every nontrivial finite abelian group G with $5G = 0$, every $\delta > 0$, and $0 < u/\delta < \eta$,*

$$\begin{aligned} \mathcal{M}_{\delta,G}(u) = & \delta^8 + 6\sqrt{2}\delta^4 u^4 + \frac{2^{3/4}}{3}\delta^2 u^6 + \frac{20}{9}u^8 \\ & + \frac{869}{216 \cdot 2^{3/4}} \frac{u^{10}}{\delta^2} + \frac{65759}{5832 \cdot 2^{3/2}} \frac{u^{12}}{\delta^4} + E_\delta(u), \end{aligned} \quad (\text{III.85.2})$$

where $|E_\delta(u)| \leq Cu^{14}/\delta^6$. More generally, $\delta^{-8}\mathcal{M}_{\delta,G}(\delta t)$ is represented near zero by a convergent power series in t^2 . The constants and the neighborhood are independent of the ambient dimension. All sufficiently small maximizers on the five-point group are obtained from one real analytic family by translations, reflection, and the choice of a generating character.

[write] **6 October 2026; sources 34–39.** *Answers source 21's question on further terms on the five-point group* (item III.27): $c_{10} = 869/(216 \cdot 2^{3/4}) = 2.39218$, $c_{12} = 65759/(5832 \cdot 2^{3/2}) = 3.98651$ (decimals recomputed at the write). The intake confirmed them on $\mathbb{Z}/5\mathbb{Z}$ by direct optimization of the cube average over two Fourier coefficients: $(M - 1 - 6\sqrt{2}u^4 - (2^{3/4}/3)u^6 - (20/9)u^8)/u^{10} = 2.4079, 2.4121, 2.4248, 2.4412, 2.4610$ at $u = 0.05, \dots, 0.13$ against $c_{10} + c_{12}u^2 = 2.4021, \dots, 2.4596$ (evidence, not proof).

The analytic neighborhood in this theorem may be smaller than the explicit range $u/\delta \leq 10^{-5}$ of exact quotient reduction. The asymptotic calculation does not require a numerical lower bound for its radius of convergence.

III.85.1 A complete Fourier certificate

Identify the dual of $\mathbb{Z}/5\mathbb{Z}$ with $\mathbb{Z}/5\mathbb{Z}$. After choosing a generator so that a largest Fourier pair is $\{1, -1\}$, write

$$a_0 = \delta, \quad a_1 = \rho e^{i\alpha}, \quad a_2 = b e^{i\beta}, \quad a_{-r} = \overline{a_r}, \quad \rho \geq b \geq 0.$$

For $b > 0$ define

$$\theta = 2\alpha - \beta, \quad \phi = 3\alpha + \beta, \quad \psi = \phi - \theta = \alpha + 2\beta.$$

If $b = 0$, any choice of the unused phase gives the same formulas. For these five Fourier coefficients set

$$T_\delta(s, t) = \sum_{r \in \mathbb{Z}/5\mathbb{Z}} a_r \overline{a_{r+s}} \overline{a_{r+t}} a_{r+s+t}.$$

The sum-of-squares identity in Lemma III.83.1 gives $\mathcal{C}_3(\delta + f) = \sum_{s,t} |T_\delta(s, t)|^2$. Realness of the underlying function makes T_δ real and invariant under independent sign changes of its two arguments and their interchange. The following six rows therefore account for all 25 slots.

Representative	Multiplicity	Value of $T_\delta(s, t)$
(0, 0)	1	$\delta^4 + 2\rho^4 + 2b^4$
(0, 1)	4	$2\delta^2\rho^2 + 2\rho^2b^2 + b^4$
(0, 2)	4	$2\delta^2b^2 + 2\rho^2b^2 + \rho^4$
(1, 1)	4	$\delta^2\rho^2 + 2\delta\rho^2b\cos\theta + 2\rho b^3\cos(2\theta - \phi)$
(2, 2)	4	$\delta^2b^2 + 2\delta\rho b^2\cos\psi + 2\rho^3b\cos\phi$
(1, 2)	8	$\rho^2b^2 + 2\delta\rho^2b\cos\theta + 2\delta\rho b^2\cos\psi$

Each row follows by substituting $r = 0, 1, 2, -2, -1$ in the defining sum; thus the table is a finite algebraic certificate, including arbitrary Fourier phases. In particular the centered norm is exactly

$$N(\rho, b, \theta, \phi) := \|f\|_{U^3}^8 = 8\rho^8 + 16(1 + \cos^2\phi)\rho^6b^2 + 48\rho^4b^4 + 16(1 + \cos^2(2\theta - \phi))\rho^2b^6 + 8b^8. \quad (\text{III.85.3})$$

Squaring all six rows gives the full objective

$$\begin{aligned} \mathcal{C}_3(\delta + f) = & \delta^8 + 24\delta^4(\rho^4 + b^4) \\ & + 16\delta^3(\rho^4b\cos\theta + \rho b^4\cos\psi) \\ & + 24\delta^2\left((3 + \cos 2\theta)\rho^4b^2 + (3 + \cos 2\psi)\rho^2b^4 \right. \\ & \quad \left. + 2\rho^3b^3(\cos\phi + \cos(2\theta - \phi))\right) \\ & + 16\delta\left(\rho^4b^3(3\cos\theta + \cos(2\phi - \theta)) \right. \\ & \quad \left. + \rho^3b^4(3\cos\psi + \cos(3\theta - \phi))\right) + N(\rho, b, \theta, \phi). \end{aligned} \quad (\text{III.85.4})$$

All optimization below uses these exact expressions. In particular, the phases are not fixed in advance.

III.85.2 Localization of every maximum at the secondary scale

By (III.85.1), take $\delta = 1$ and put $t = u$. The norm sphere of centered real functions on the five-point group is compact, so the maximum exists. Since $\rho \geq b$, the nonnegative terms in (III.85.3) imply

$$8\rho^8 \leq t^8 \leq 128\rho^8. \quad (\text{III.85.5})$$

Thus $\rho \asymp t$ uniformly in the phases, and both amplitudes are $O(t)$.

Lemma III.85.2 (Uniform secondary localization). *Every five-point maximizer has $b = O(t^2)$ as $t \downarrow 0$, with an absolute implied constant.*

Proof. Put $S = 2(\rho^4 + b^4)$. Equation (III.85.3) gives $t^8 - 2S^2 \geq 16\rho^6 b^2$. Therefore

$$6\sqrt{2}t^4 - 12S = \frac{6(t^8 - 2S^2)}{t^4/\sqrt{2} + S} \geq ct^2 b^2$$

for an absolute $c > 0$, using (III.85.5) and $S \leq t^4/\sqrt{2}$. The fifth-, sixth-, and seventh-degree parts of (III.85.4) have moduli at most, respectively, $Ct^4 b$, $Ct^4 b^2$, and $Ct^4 b^3$. For example $\rho b^4 \leq \rho^4 b$, $\rho^2 b^4 \leq \rho^4 b^2$, and $\rho^3 b^3 \leq \rho^4 b^2$, because $b \leq \rho$. The constants are absolute because the displayed polynomial has finitely many terms with bounded trigonometric factors.

The pure-pair function $b = 0$ has value $1 + 6\sqrt{2}t^4 + t^8$. For a maximizer, comparison with this function and absorption of $Ct^4 b^2 + Ct^4 b^3$ into $ct^2 b^2/2$, for small t , give

$$0 \leq C_3(1 + f) - 1 - 6\sqrt{2}t^4 - t^8 \leq -\frac{c}{2}t^2 b^2 + Ct^4 b.$$

If $b > 0$, division by $t^2 b$ yields $b \leq (2C/c)t^2$; the claim is immediate if $b = 0$. $\square$

This lemma supplies the compactness needed for a perturbative maximum. It rules out functions whose second harmonic leaves every bounded region at the scale relevant to the next coefficient.

III.85.3 Analytic optimization with both phases retained

Set

$$q_0 = 2^{-3/4}, \quad b = kt^2, \quad \rho^2 = t^2 R.$$

After dividing the norm equation by t^8 , one obtains

$$\begin{aligned} 1 &= 8R^4 + 16(1 + \cos^2 \phi)R^3 k^2 t^2 + 48R^2 k^4 t^4 \\ &\quad + 16(1 + \cos^2(2\theta - \phi))Rk^6 t^6 + 8k^8 t^8. \end{aligned} \quad (\text{III.85.6})$$

At $t = 0$, the positive solution is $R = q_0$, and the derivative of the right side with respect to R is $32q_0^3 > 0$. The implicit function theorem gives a positive real analytic solution $R(t, k, \theta, \phi)$ near zero. It is analytic in t^2 , uniformly for bounded k and all phases: this follows either from a finite cover of the phase torus or directly from the uniform nonzero derivative at $R = q_0$.

Define

$$H(t, k, \theta, \phi) = \frac{C_3(1 + f) - 1 - 6\sqrt{2}t^4}{t^6}, \quad (\text{III.85.7})$$

with $\rho = t\sqrt{R}$ for positive t . The apparent singularity at $t = 0$ is removable, as substitution in (III.85.4) shows. The resulting expression has a real analytic continuation to signed t using this same formula for ρ . Its expansion begins

$$H = H_0 + t^2 H_2 + t^3 H_3 + O(t^4), \quad (\text{III.85.8})$$

where, writing $A = 1 + \cos^2 \phi$,

$$H_0 = -24q_0 A k^2 + 16q_0^2 k \cos \theta, \quad (\text{III.85.9})$$

$$H_2 = 1 + 24(A^2 - 2)k^4 - 16q_0 A k^3 \cos \theta + 24(3 + \cos 2\theta)q_0^2 k^2, \quad (\text{III.85.10})$$

$$H_3 = 16\sqrt{q_0} k^4 \cos(\phi - \theta) + 48q_0^{3/2} k^3 (\cos(2\theta - \phi) + \cos \phi). \quad (\text{III.85.11})$$

For example the first two corrections in the norm solution are

$$R = q_0 - \frac{A}{2}k^2t^2 + \frac{3(A^2 - 4)}{8q_0}k^4t^4 + O(t^6),$$

which already verify (III.85.9) and (III.85.10).

Proposition III.85.3 (All local maximizing branches). *For sufficiently small positive t , every global maximum is on an analytic critical branch converging, in the coordinates (k, θ, ϕ) , to one of*

$$y_+ = (q_0/3, 0, \pi/2), \quad y_- = (q_0/3, 0, 3\pi/2), \quad (\text{III.85.12})$$

with phases understood modulo 2π . There is exactly one critical branch near each class. They have equal objective values.

Proof. For $k \geq 0$,

$$H_0 \leq -24q_0k^2 + 16q_0^3k \leq \frac{8q_0^3}{3} = \frac{2^{3/4}}{3}.$$

Equality in the second inequality requires $k = q_0/3 > 0$. Equality in the first then forces $\cos \theta = 1$ and $\cos \phi = 0$, yielding exactly (III.85.12). Lemma III.85.2 confines all maximizers to a fixed compact set of k and the phase torus. On that set $H \rightarrow H_0$ uniformly. Each of the points in (III.85.12), together with its analytic norm solution, is an admissible competitor for small positive t . Compactness therefore forces every sequence of maximizers with $t \downarrow 0$ to approach one of these two classes.

At either class the Hessian of H_0 , in coordinates (k, θ, ϕ) , is

$$\text{diag} \left(-48q_0, -\frac{16q_0^3}{3}, -\frac{16q_0^3}{3} \right). \quad (\text{III.85.13})$$

It is negative definite. Apply the implicit function theorem to $\nabla_{k,\theta,\phi}H = 0$ at each point. This gives one analytic critical branch in a neighborhood of each class, and its Hessian remains negative definite. The preceding compactness argument, applied to these smaller neighborhoods, shows that these branches exhaust the global maxima. Their amplitude constraints are inactive: k remains positive and bounded, while $b/\rho = O(t)$. Conjugating the Fourier coefficients sends (θ, ϕ) to $(-\theta, -\phi)$ and interchanges the branches without changing the objective, so their values agree. $\square$

The map $(\alpha, \beta) \mapsto (\theta, \phi)$ has determinant 5. Its five lifts on the two-dimensional phase torus correspond exactly to translations of the five-point function. Choosing the other largest-pair generator is a group automorphism, and conjugation of the Fourier coefficients is reflection. This proves the symmetry assertion in Theorem III.85.1.

III.85.4 Evenness and a recursion for all further coefficients

Use the branch near y_+ and view $\phi - \pi/2$ as a local real coordinate. The exact normalized equations have the symmetry

$$(t, k, \theta, \phi) \mapsto (-t, k, -\theta, \pi - \phi). \quad (\text{III.85.14})$$

To check it, note that $\cos \theta$, $\cos 2\theta$, $\cos 2\psi$, and $\cos(2\phi - \theta)$ are unchanged, whereas $\cos \psi$, $\cos \phi$, $\cos(2\theta - \phi)$, and $\cos(3\theta - \phi)$ change sign. The latter factors occur precisely with the odd powers of signed t in the normalized objective. Both squared cosine factors in (III.85.6) are unchanged.

The involution fixes y_+ . Uniqueness of the analytic critical branch thus implies that $k(t)$ is even, $\theta(t)$ is odd, $\phi(t) - \pi/2$ is odd, and its objective value is even. Consequently the extremal profile is analytic in t^2 , rather than merely possessing a finite asymptotic expansion.

This gives an effective recursion. First solve (III.85.6) coefficient by coefficient for R ; the coefficient to be solved is multiplied by $32q_0^3$. Then expand $\nabla H = 0$. At each successive order, the new coefficient vector of (k, θ, ϕ) is multiplied by the fixed invertible matrix (III.85.13); all other terms are already known from lower orders. Every coefficient is therefore determined by finite algebraic operations. The analytic implicit function theorem proves convergence of the recursively obtained series.

III.85.5 Fourth, sixth, eighth, and tenth orders

At the limiting phases $\theta = 0$, $\phi = \pi/2$, put $z = k^2 t^2 / q_0$ and write $R = q_0 X(z)$. The exact norm becomes

$$X^4 + 2zX^3 + 6z^2X^2 + 2z^3X + z^4 = 1, \quad X(0) = 1. \quad (\text{III.85.15})$$

Coefficient comparison gives

$$X = 1 - \frac{1}{2}z - \frac{9}{8}z^2 + \frac{3}{4}z^3 + \frac{39}{128}z^4 + O(z^5). \quad (\text{III.85.16})$$

At these phases the objective simplifies to

$$\mathcal{C}_3(1 + f) = 1 + 24(\rho^4 + b^4) + 16\rho^4 b + 96\rho^4 b^2 + 48\rho^2 b^4 + 32\rho^4 b^3 + t^8. \quad (\text{III.85.17})$$

Its first coefficient beyond the leading quartic term is the maximum $8q_0^3/3 = 2^{3/4}/3$ of H_0 . At these fixed phases the next coefficients of H are

$$H_2(k) = 1 - 24k^4 - 16q_0k^3 + 96q_0^2k^2, \quad (\text{III.85.18})$$

$$H_4(k) = \frac{63k^6}{q_0} - 32k^5 - 48q_0k^4 + 32q_0^2k^3, \quad (\text{III.85.19})$$

$$H_6(k) = \frac{27k^8}{q_0^2} + \frac{42k^7}{q_0} - 216k^6 - 32q_0k^5. \quad (\text{III.85.20})$$

Here $H_j(k)$ denotes the coefficient of t^j after the phases are set to their limiting values; the arbitrary-phase functions H_2, H_3 remain those in (III.85.10)–(III.85.11).

Let $k_0 = q_0/3$. Since $q_0^4 = 1/8$,

$$H_2(k_0) = 1 - \frac{1}{27} - \frac{2}{27} + \frac{4}{3} = \frac{20}{9}. \quad (\text{III.85.21})$$

The analytic maximizing coordinates have no term of order t , because (III.85.8) has none and the Hessian is invertible. Equation (III.85.10) shows that its phase gradient vanishes at y_+ . Thus the displacement at order t^2 occurs only in the amplitude coordinate:

$$k(t) = k_0 + k_2 t^2 + O(t^4), \quad k_2 = -\frac{H'_2(k_0)}{H''_0(k_0)} = \frac{31q_0^2}{27}, \quad (\text{III.85.22})$$

where

$$H'_2(k_0) = \frac{496q_0^3}{9}, \quad H''_0(k_0) = -48q_0.$$

The optimized coefficient at order t^4 in H , hence at order t^{10} in the original objective, is

$$\begin{aligned} K_{10} &= H_4(k_0) + \frac{(H'_2(k_0))^2}{96q_0} \\ &= \frac{133q_0^5}{243} + \frac{7688q_0^5}{243} = \frac{869q_0}{216} = \frac{869}{216 \cdot 2^{3/4}}. \end{aligned} \quad (\text{III.85.23})$$

The phase displacement begins at order t^3 , so it cannot affect this coefficient. The next subsection computes its first effect.

III.85.6 The first phase adjustment and the twelfth-order coefficient

At y_+ , differentiation of (III.85.11) gives

$$\nabla_{\theta,\phi} H_3(y_+) = \frac{304q_0^{9/2}}{81}(1, -1), \quad \partial_k H_3(y_+) = 0. \quad (\text{III.85.24})$$

The phase part of the Hessian is $-(16q_0^3/3)I$. The order- t^3 critical-point equations therefore yield

$$\begin{aligned} \theta(t) &= d_3 t^3 + O(t^5), \\ \phi(t) &= \frac{\pi}{2} - d_3 t^3 + O(t^5), \quad d_3 = \frac{19q_0^{3/2}}{27}. \end{aligned} \quad (\text{III.85.25})$$

The absence of fourth-order phase terms follows from (III.85.14). In the original phases, one lift has

$$\alpha(t) = \frac{\pi}{10} + O(t^5), \quad \beta(t) = \frac{\pi}{5} - \frac{19q_0^{3/2}}{27} t^3 + O(t^5).$$

Thus keeping both phases permanently at their limiting values is already insufficient for the twelfth-order extremal coefficient.

We give the coefficient calculation with its optimization terms retained. Write $y(t) = y_+ + t^2 y_2 + t^3 y_3 + t^4 y_4 + \dots$, where $y_2 = (k_2, 0, 0)$ and $y_3 = (0, d_3, -d_3)$. The coefficient of t^6 in $H(t, y(t))$ is

$$\begin{aligned} K_{12} &= H_6(k_0) + k_2 H'_4(k_0) + \frac{1}{2} k_2^2 H''_2(k_0) \\ &\quad - \frac{1}{2} (\nabla_{\theta,\phi} H_3(y_+))^T (\nabla_{\theta,\phi}^2 H_0(y_+))^{-1} \nabla_{\theta,\phi} H_3(y_+). \end{aligned} \quad (\text{III.85.26})$$

To justify this formula, the terms containing the unknown y_4 cancel because $\nabla^2 H_0 y_2 + \nabla H_2 = 0$. The cubic derivative of H_0 in the direction y_2 is zero, since its restriction to the amplitude coordinate is quadratic. The terms involving y_3 complete the displayed quadratic form, using $\nabla^2 H_0 y_3 + \nabla H_3 = 0$. No term involving H_5 contributes at this order: the coordinate displacement starts at t^2 , and H_5 has prefactor t^5 . These observations account for every possible term through t^6 .

The needed fixed-phase values are

$$H_6(k_0) = -\frac{295q_0^6}{729}, \quad H'_4(k_0) = \frac{254q_0^4}{81}, \quad H''_2(k_0) = 128q_0^2.$$

The amplitude contribution in (III.85.26) and the phase contribution are, respectively,

$$k_2 H'_4(k_0) + \frac{1}{2} k_2^2 H''_2(k_0) = \frac{192386q_0^6}{2187}, \quad \frac{5776q_0^6}{2187}.$$

Consequently

$$K_{12} = \frac{65759q_0^6}{729} = \frac{65759}{5832 \cdot 2^{3/2}}. \quad (\text{III.85.27})$$

The phase contribution is strictly positive. It quantifies the first improvement obtained by optimizing the phases beyond their limiting values.

Completion of Theorem III.85.1. Proposition III.85.3 identifies all maxima near zero with the analytic branches just computed. Symmetry (III.85.14) makes their common value even. Equations (III.85.21), (III.85.23), and (III.85.27) therefore give an error $O(t^{14})$ after restoring the constant and quartic terms in (III.85.7). Its constant is absolute, because this is an analytic function on a fixed finite-dimensional neighborhood. Scale by (III.85.1) and apply Theorem III.84.1. This proves the claimed remainder and its independence of the ambient dimension. $\square$

III.85.7 Sharp realization by bounded weights and exact verification

Corollary III.85.4 (Weighted extremizers). *Let $0 < \delta < 1$ and let G be nontrivial with $5G = 0$. If $0 < u \leq 10^{-5}\delta$ and $4u \leq \min\{\delta, 1 - \delta\}$, then $\mathcal{M}_{\delta,G}(u)$ is attained by a weight $W : G \rightarrow [0, 1]$ of mean δ . The same maximum is obtained when the centered norm constraint is $\|W - \delta\|_{U^3} \leq u$. In particular, all coefficients in Theorem III.85.1 are sharp for bounded weights as well as for unrestricted real functions.*

Proof. Every five-point centered function has at most four nonzero Fourier coefficients. Each coefficient has modulus at most its U^2 norm, and hence at most its U^3 norm. Thus $\|f\|_\infty \leq 4u$ for a centered five-point function of norm u . The hypotheses place $\delta + f$ in $[0, 1]$. Pullback along a nontrivial character of G preserves the mean, the norm, and the cube count. The unrestricted upper bound then proves equality for weights. The norm-ball assertion follows from the strict monotonicity in Theorem III.84.1. $\square$

Corollary III.85.5 (An indicator version of the cube-count bound). *Let G be nontrivial with $5G = 0$, let $A \subseteq G$ have density $\delta > 0$, and suppose*

$$\|\mathbf{1}_A - \delta\|_{U^3}^8 \leq \alpha, \quad \alpha^{1/8} \leq 10^{-5}\delta.$$

The number of ordered cube parameters $(x, h_1, h_2, h_3) \in G^4$ whose eight vertices all belong to A is at most

$$|G|^4 \mathcal{M}_{\delta,5}(\alpha^{1/8}). \quad (\text{III.85.28})$$

Here $\mathcal{M}_{\delta,5}(0) = \delta^8$ by continuous extension. Repeated vertices and zero increments are included. As $\alpha^{1/8}/\delta \rightarrow 0$, its normalized upper bound is

$$\begin{aligned} C_3(\mathbf{1}_A) \leq & \delta^8 + 6\sqrt{2}\delta^4\alpha^{1/2} + \frac{2^{3/4}}{3}\delta^2\alpha^{3/4} + \frac{20}{9}\alpha \\ & + \frac{869}{216 \cdot 2^{3/4}}\delta^{-2}\alpha^{5/4} + \frac{65759}{5832 \cdot 2^{3/2}}\delta^{-4}\alpha^{3/2} + C\delta^{-6}\alpha^{7/4}, \end{aligned} \quad (\text{III.85.29})$$

for an absolute C , on a sufficiently small absolute range of $\alpha^{1/8}/\delta$.

Proof. Apply the norm-ball form of Theorem III.84.1 to $\mathbf{1}_A - \delta$, and multiply its normalized cube count by $|G|^4$. If $\alpha = 0$, the centered function vanishes and the assertion is immediate. Otherwise Theorem III.85.1, with $u = \alpha^{1/8}$, gives the displayed powers and the positive remainder bound. $\square$

Gowers' Lemma 3.10 gives the normalized bound $(\delta + \alpha^{1/8})^8$ in this notation [1, Lemma 3.10]. Its first error is of order $\alpha^{1/8}$ at fixed positive density, whereas (III.85.29) begins at $\alpha^{1/2}$. The sharp leading coefficient $6\sqrt{2}$ was already obtained in the preceding project reports [169, 170]; the additions here are exact quotient reduction and the dimension-uniform higher coefficients. The corollary is an upper bound for indicators; the exact extremizers and sharp realizations proved above are weights.

The program `code/verify_twelfth_order.py` checks (III.85.16), all the displayed H_j coefficients, the Hessian and phase displacements, and both optimized higher coefficients by exact symbolic arithmetic. It records the checks in `data/twelfth_order_checks.json`. A separate high-precision stationary-branch calculation in the same program compares the twelfth-order residual at decreasing positive values of t . That numerical calculation checks the analytic branch; it is not used to establish global optimality or the size of its neighborhood.

III.86 Cyclic reduction for every fixed odd exponent

Source 38, delivered Section 5.

The mechanism behind exact reduction is not exclusive to characteristic five. The finite size of the cyclic frequency group generated by a dominant character is what bounds the potentially negative cross terms. The following version keeps the dependence on that size explicit.

Theorem III.86.1 (Exact cyclic reduction at bounded odd exponent). *Let $m \geq 3$ be odd, let G be a finite abelian group of exponent exactly m , and put $c_m = 10^{-8}m^{-8}$. For every $\delta > 0$,*

$$\mathcal{M}_{\delta,G}(u) = \mathcal{M}_{\delta,\mathbb{Z}/m\mathbb{Z}}(u) \quad (0 < u \leq c_m\delta). \quad (\text{III.86.1})$$

Every maximizing function on G factors through a cyclic quotient of odd order dividing m . The same identity of maxima holds with the norm constraint replaced by $\|f\|_{U^3} \leq u$. For $0 < \delta < 1$, these unrestricted maxima are attainable by weights in $[0, 1]$ if also $mu \leq \min\{\delta, 1 - \delta\}$.

The threshold is a sufficient bound; neither its constant nor its power of m is claimed optimal. In particular, this theorem does not supply a threshold independent of all odd exponents.

Proof. Normalize $f = uF$, $\|F\|_{U^3} = 1$, and $t = u/\delta$. We first verify the pure-pair benchmark used in localization. For a real function with Fourier support $\{\xi, -\xi\}$ and coefficients $\rho e^{i\alpha}$ and its conjugate, the centered norm is $8\rho^8$ for every nonzero odd-order character. Indeed a two-point subset of an odd group has no parallelogram with both increments nonzero, so only the axes of the Fourier-square formula contribute. If the character has order greater than three, the fifth-, sixth-, and seventh-degree terms vanish. For the sixth and seventh degrees this follows from the statistic H : no three frequencies from this pair sum to zero.

If the character has order three, the terms are instead

$$B_5 = 2\rho^5 \cos 3\alpha, \quad C_6 = (24 + 8 \cos 6\alpha)\rho^6, \quad C_7 = 0.$$

These identities follow directly from the H, J, K and T, L formulas: here $H = 2\rho^6$, $J = 0$, $K = 2\rho^6 \cos 6\alpha$, and the nonzero slots of L are disjoint from those of T . Choosing $\alpha = 0$ makes both additional terms positive. Thus every nontrivial odd-order group has a normalized pure-pair competitor of value at least $1 + 6\sqrt{2}t^4 + t^8$.

For any maximizing candidate put $S = \sum_r |\hat{F}(r)|^4$ and $S_0 = 2^{-1/2}$. Lemma III.83.3, cube Cauchy–Schwarz, and the benchmark imply $0 \leq S_0 - S \leq 4t$. Choose a largest Fourier pair, generated by ξ , and put $D_{\text{pair}} = \sum_{r \notin \{\xi, -\xi\}} |\hat{F}(r)|^4$. Lemma III.83.2 gives

$$1 \geq 2S^2 + 3SD_{\text{pair}}, \quad D_{\text{pair}} \leq 8t$$

in the present range. Let $L = \langle \xi \rangle$, of size at most m , let g be the Fourier projection to L , and write $h = F - g$. As in the exact reduction proof, put

$$S_g = \sum_{r \in L} |\hat{F}(r)|^4, \quad D = \sum_{r \notin L} |\hat{F}(r)|^4, \quad v = \|g\|_{U^3}, \quad \Delta = 1 - v^8.$$

Then $S_g \geq S_0 - 12t$.

Let p denote the Fourier projection of g onto the principal pair. For $s, t \in L$ with $s, t \neq 0$, the preceding two-point observation gives $T_p(s, t) = 0$. Telescope the four factors of $T_g - T_p$ and apply ℓ^4 Hölder to each of the four terms. Each unchanged factor has norm at most one, whereas the changed factor has norm at most $D_{\text{pair}}^{1/4}$. Therefore

$$|T_g(s, t)| \leq 4D_{\text{pair}}^{1/4}, \quad R := \sum_{\substack{s, t \in L \\ s, t \neq 0}} |T_g(s, t)| \leq 4m^2(8t)^{1/4}.$$

For $t \leq c_m$, this is at most $4 \cdot 8^{1/4}/100 < 1/3$. Also $S_g \geq 2/3$. The energy argument proving (III.84.1) uses only odd order and a subgroup L , before its final five-point bound on R is inserted. It therefore gives here

$$\Delta \geq (10S_g - 2R)D + 2D^2 \geq 6D, \quad 0 \leq \Delta \leq 1/9. \quad (\text{III.86.2})$$

The last bound follows from $v^8 \geq 2S_g^2 \geq 8/9$. The subgroup-projection estimate (III.84.2) still gives $\|h\|_{U^3}^8 \leq 256\Delta$.

We spell out the dependence of the higher-degree estimates on m . Write H_F, H_g for the nonnegative triple statistics of the two functions. Cauchy–Schwarz gives

$$A_g := \sum_{r \in L} |\hat{F}(r)|^2 \leq \sqrt{|L|S_g} \leq \sqrt{m}.$$

Exactly as in the proof of (III.84.3),

$$\begin{aligned} H_F - H_g &\leq 3A_g D + \sqrt{D(128\Delta + D^2/2)} \\ &\leq C_H \Delta, \quad C_H = 5 + \frac{\sqrt{m}}{2} \leq 3m. \end{aligned}$$

Thus $|C_6(F) - C_6(g)| \leq 28C_H \Delta$. For the seventh-degree difference the part on L^2 is at most $8|L|^2 D \leq (4/3)m^2 \Delta$, because each value of $L_g(s, t)$ has modulus at most four. On the complement, the two squared norms in the Cauchy–Schwarz estimate are at most $(4/3)\Delta$ and $16C_H \Delta$. Therefore

$$|C_7(F) - C_7(g)| \leq \left(\frac{4}{3}m^2 + \frac{16}{\sqrt{3}}\sqrt{C_H} \right) \Delta.$$

This is the same linear estimate as (III.84.4), now with its dependence on m retained.

Multiplication by two preserves both L and its complement, because the group has odd order. Consequently $B_5(F) - B_5(g) = B_5(h)$ and its modulus is at most $D^{5/4} \leq \Delta/6$. Set $g_* = g/v$. For

$k = 5, 6, 7$, the rescaling estimate $|v^{-k} - 1| \leq 4\Delta$ remains valid. Combining it with the preceding bounds and cube Cauchy–Schwarz yields

$$\begin{aligned} |B_5(F) - B_5(g_*)| &\leq 5\Delta, \\ |C_6(F) - C_6(g_*)| &\leq (28C_H + 112)\Delta \leq 122m\Delta, \\ |C_7(F) - C_7(g_*)| &\leq 10m^2\Delta. \end{aligned}$$

For the last simplification, $C_H \leq 3m$ and $m \geq 3$ give $16\sqrt{C_H/3} \leq 16\sqrt{m} \leq 10m$; the remaining terms are $(4/3)m^2 + 32$, whose sum with $10m$ is less than $10m^2$.

The fourth-degree improvement is at least $2t^4\Delta$, by (III.86.2), and the eighth-degree terms cancel because both functions have norm one. Hence

$$C_3(1 + tF) - C_3(1 + tg_*) \leq t^4\Delta(-2 + 40t + 122mt^2 + 10m^2t^3) \leq -t^4\Delta.$$

The coefficient in parentheses is strictly less than -1 throughout $t \leq 10^{-8}m^{-8}$; substitution bounds the three positive terms by $40 \cdot 10^{-8}m^{-8}$, $122 \cdot 10^{-16}m^{-15}$, and $10 \cdot 10^{-24}m^{-22}$, respectively, with sum less than one. If $D > 0$, the improvement is strict. Every maximizer consequently has Fourier support in a cyclic group L of order $d \mid m$.

The dual of G has an element of order m , so G has a cyclic quotient of that order. Conversely every function on a cyclic group of order $d \mid m$ pulls back to $\mathbb{Z}/m\mathbb{Z}$ with its mean, norm, and objective preserved. These two pullback statements prove (III.86.1).

For the norm-ball assertion, let F be normalized and maximizing at centered norm $0 < r \leq c_m$. The localization above gives $S \geq 2/3$. For $r \leq s \leq c_m$, the centered polynomial satisfies

$$\frac{d}{ds}C_3(1 + sF) \geq s^3(48S - 40s - 168s^2 - 56s^3) > 0.$$

Scaling a positive lower-norm maximizer therefore proves strict monotonicity throughout the range. The norm-zero value is strictly smaller by the pure-pair benchmark. The ball maximum is attained at its boundary. Finally, a centered function on $\mathbb{Z}/m\mathbb{Z}$ has at most $m - 1$ nonzero Fourier coefficients, each of modulus at most u . The additional bound $mu \leq \min\{\delta, 1 - \delta\}$ puts its maximizing weight in $[0, 1]$, and pullback concludes the proof. $\square$

Remark III.86.2 (Extremal values and the order of individual quotients). The equality of extremal values with $\mathbb{Z}/m\mathbb{Z}$ does not say that every maximizing function itself factors through an order- m quotient. A lower-order character need not lie in a cyclic subgroup of order m : in the dual of $\mathbb{Z}/9\mathbb{Z} \times \mathbb{Z}/3\mathbb{Z}$, for example, the character represented by $(0, 1)$ is not contained in any cyclic subgroup of order nine. The theorem gives cyclic dependence of each extremizer, while the extremal value is compared by pullbacks between the resulting cyclic groups.

III.87 Source 38's Part III: six-term norm-extension defects [write]

Added in the seventh write, 6 October 2026.

The next section is source 38's Section 7, its delivered Part III. *Logically it follows source 22's chapter on polynomial colourings* (Section III.60), whose cubic-core defect and zero-defect classification it extends, and the norm-colouring material of sources 10 and 14; it is printed here because no number printed earlier may move. For a layered linear–cubic–quintic norm colouring plus an

arbitrary function of independent coordinates it counts the symmetric six-term progressions exactly; the excess over the surjectivity count is governed by centered third and fifth differences and a quadratic-character statistic (Theorem III.88.2). The layered construction is Shi and Dong’s; no better colouring exponent is claimed, and the quintic four-term dimension question stays open (item III.19). Nothing in this section is formalized.

III.88 An exact six-term extension defect and its quadratic-character term

Source 38, delivered Section 7.

*Delivered part heading of source 38: **Exact six-term norm-extension defects***

Gowers’ Conjecture 4.2 proposed a random-density lower bound for even-length progressions under a weaker uniformity hypothesis [1]. Symmetric-progression colorings are relevant to the limits of such a principle. In particular, Shi and Dong use layered odd-degree field norms to construct colorings avoiding symmetric progressions, with applications to counterexamples to the conjecture [33]. The question studied here is an exact counting problem inside that construction method: what happens after a function of new, independent coordinates is added to a complete norm core?

The preceding degree-stratification report [173] proves that such an addition cannot preserve avoidance, and gives an exact third-difference formula for the linear–cubic core. It asks for exact formulas beyond that case. We answer the next case. The six-term core contains both a cubic and a quintic norm; after one averages over its centers, their interaction is governed by the value distribution of a nondegenerate quadratic form on a five-dimensional extension field. The resulting formula applies to an arbitrary added function, without a polynomial-degree assumption.

III.88.1 Fields, extension degrees, and norm fibers

Fix a prime power $q = p^r$ of characteristic $p > 5$. Let N_j denote the norm from $\mathbb{F}_{q^j}$ to $\mathbb{F}_q$. The subscripts 3 and 5 below are *extension degrees over $\mathbb{F}_q$* ; neither is the characteristic. The base field need not be prime. The assumption concerns p , not merely q : it guarantees six distinct progression points in every nonzero direction, the invertibility of 120, and the nonvanishing trace $\mathrm{Tr}_{\mathbb{F}_{q^5}/\mathbb{F}_q}(1) = 5$.

Lemma III.88.1 (Norm presentation and fibers). *After any linear identification $\mathbb{F}_{q^j} \cong \mathbb{F}_q^j$, the norm N_j is represented by a homogeneous polynomial of total degree j . Its zero fiber is $\{0\}$, and each nonzero scalar has exactly*

$$L_j = \frac{q^j - 1}{q - 1}$$

preimages.

Proof. In a basis $b_1, \dots, b_j$, write $x = \sum_i X_i b_i$. Then

$$N_j(x) = \prod_{s=0}^{j-1} \left(\sum_i X_i b_i^{q^s} \right).$$

This is a product of j linear forms. Frobenius permutes the factors, so its coefficients belong to $\mathbb{F}_q$. The product vanishes precisely when $x = 0$. On $\mathbb{F}_{q^j}^*$ it is the homomorphism $x \mapsto x^{L_j}$. The

multiplicative group is cyclic of order $q^j - 1$, and the norm of a generator has order $q - 1$. Thus the map is onto $\mathbb{F}_q^*$ and has kernel of size L_j , which gives every nonzero fiber size. $\square$

In particular, put

$$L_3 = q^2 + q + 1, \quad L_5 = q^4 + q^3 + q^2 + q + 1, \quad \chi : \mathbb{F}_q \longrightarrow \{-1, 0, 1\},$$

where χ is the quadratic character extended by $\chi(0) = 0$. For an integer $e \geq 0$ and an arbitrary function $R : \mathbb{F}_q^e \rightarrow \mathbb{F}_q$, consider

$$P(a, u, v, w) = a + N_3(u) + N_5(v) + R(w).$$

Here $(a, u, v, w) \in \mathbb{F}_q \times \mathbb{F}_{q^3} \times \mathbb{F}_{q^5} \times \mathbb{F}_q^e$, so its domain has dimension $9 + e$ over $\mathbb{F}_q$. Write $A_6^{\text{sym}}(P)$ for the number of ordered center–direction pairs (z, h) with $h \neq 0$ such that

$$P(z + h/2) = P(z - h/2), \quad P(z + 3h/2) = P(z - 3h/2), \quad P(z + 5h/2) = P(z - 5h/2).$$

This counts ordered nontrivial six-term arithmetic progressions with the three stated symmetric color equalities. The direction is the common difference, and no quotient by sign or by rescaling is taken.

For $b, g \in \mathbb{F}_q^e$, define the centered odd differences

$$\begin{aligned} D_3(b, g) &= \Delta_g^3 R(b - 3g/2), \\ D_5(b, g) &= \Delta_g^5 R(b - 5g/2), \end{aligned}$$

where $\Delta_g R(x) = R(x + g) - R(x)$. Define three finite statistics

$$\begin{aligned} T_0 &= \#\{(b, g) : D_5(b, g) = 0, D_3(b, g) \neq 0\}, \\ T_1 &= \#\{(b, g) : D_5(b, g) \neq 0\}, \\ S &= \sum_{\substack{b, g \\ D_5(b, g) \neq 0}} \chi\left(\frac{40D_3(b, g)}{D_5(b, g)} - 5\right). \end{aligned}$$

All quotients are taken in $\mathbb{F}_q$. The integers appearing in the character argument are reduced in the same field.

Theorem III.88.2 (Exact six-term defect). *With the notation above,*

$$\begin{aligned} A_6^{\text{sym}}(P) &= q^{9+e}(q^e - 1) + q^9(q^2 + q)T_0 \\ &\quad + q^9(q^2 L_5 - 1)T_1 - q^7(q + 1)L_5 S. \end{aligned} \tag{III.88.1}$$

In particular the general surjectivity lower bound $q^{9+e}(q^e - 1)$ is attained exactly when all pure third additive differences of R vanish. Equivalently, after regarding its domain as a vector space over $\mathbb{F}_p$, the function R has the form

$$R(x) = c + L(x) + B(x, x),$$

where L is $\mathbb{F}_p$ -linear and B is symmetric $\mathbb{F}_p$ -bilinear, with values in $\mathbb{F}_q$.

[write] **6 October 2026; sources 34–39.** *In this report.* This extends source 22’s exact cubic-core defect and zero-defect classification (Corollary III.64.4) to the linear–cubic–quintic core with an independent addition; source 22’s question on defect formulas for more odd norm layers (item III.28) is advanced for three layers. The layered construction is Shi and Dong’s, as the source says; no better colouring exponent is claimed. The exact finite checks over the fields of orders 7 and 13 were rerun on a copy at the write.

The formula gives a sharper hierarchy than the statement that every added direction has at least one matching core direction. A nonzero third difference with zero fifth difference contributes $q^2 + q$ extra core directions on average. A nonzero fifth difference contributes approximately q^6 extra core directions on average; its exact deviation has a quadratic-character sign. The conclusion concerns this specified independent extension of a norm coloring. It does not improve the unrestricted symmetric-progression coloring exponent or the global Szemerédi bound.

III.88.2 The quadratic coefficient of the quintic norm

For $x \in E = \mathbb{F}_{q^5}$, write

$$N_5(x + T) = T^5 + e_1(x)T^4 + e_2(x)T^3 + e_3(x)T^2 + e_4(x)T + e_5(x).$$

Expanding the square of the sum of the five Frobenius conjugates gives

$$e_2(x) = \frac{(\mathrm{Tr} x)^2 - \mathrm{Tr}(x^2)}{2}.$$

Lemma III.88.3 (Exact value distribution). *For every $c \in \mathbb{F}_q$,*

$$\#\{x \in E : e_2(x) = c\} = q^4 + q^2\chi(2c).$$

Proof. Choose an $\mathbb{F}_q$ -basis $b_1, \dots, b_5$ of E , and let $T = (\mathrm{Tr}(b_i b_j))_{i,j}$ be the trace-pairing matrix. The pairing is nondegenerate: for $x \neq 0$, pairing with x^{-1} gives $\mathrm{Tr}(1) = 5 \neq 0$. Its determinant is a nonzero square in $\mathbb{F}_q$. Indeed, if $M = (b_j^{q^i})_{0 \leq i \leq 4, 1 \leq j \leq 5}$, then $T = M^\top M$; moreover Frobenius acts on M by a five-cycle of its rows, so $(\det M)^q = \det M$. Thus $\det M \in \mathbb{F}_q^*$ and $\det T = (\det M)^2$.

The quadratic form e_2 is represented relative to the trace pairing by the self-adjoint operator

$$\frac{1}{2}(1 \otimes \mathrm{Tr} - I).$$

The operator in parentheses has eigenvalue 4 on the span of 1 and eigenvalue -1 on the trace-zero subspace. These subspaces are complementary because $\mathrm{Tr}(1) = 5 \neq 0$. Hence its determinant is 4, and the determinant of a matrix for e_2 is

$$2^{-5} \cdot 4 \det T = \frac{\det T}{8}.$$

In particular, e_2 is nondegenerate of rank five and its determinant has quadratic character $\chi(2)$.

For completeness, the standard odd-rank quadratic-form count follows by diagonalization and additive-character orthogonality. For a nontrivial additive character ψ , put $G = \sum_{x \in \mathbb{F}_q} \psi(x^2) = \sum_t \chi(t)\psi(t)$. The identity $G^2 = \chi(-1)q$ follows by writing the two nonzero variables in G^2 as u, ut : the sum over $u \neq 0$ of $\psi(u(1+t))$ is $q-1$ at $t = -1$ and -1 otherwise. If Q is a nondegenerate quadratic form of rank five with matrix A , then

$$\sum_x \psi(sQ(x)) = \chi(s)\chi(\det A)G^5 \quad (s \neq 0).$$

Orthogonality gives q^4 solutions when $c = 0$. When $c \neq 0$, use $\sum_{s \neq 0} \chi(s)\psi(-sc) = \chi(-c)G$ to obtain

$$\#\{x : Q(x) = c\} = q^4 + q^{-1}\chi(\det A)\chi(-c)G^6 = q^4 + q^2\chi(c \det A).$$

Applying this with $\chi(\det A) = \chi(2)$ proves the result. $\square$

III.88.3 The exact averaged reflection fibers

For fixed (b, g) , let $\beta_1, \beta_3, \beta_5$ be the unique coefficients such that

$$R(b + \alpha g) - R(b - \alpha g) = 2(\beta_1 \alpha + \beta_3 \alpha^3 + \beta_5 \alpha^5) \quad \left(\alpha = \frac{1}{2}, \frac{3}{2}, \frac{5}{2} \right).$$

The matrix is invertible because the three nonzero squared parameters are distinct when $p > 5$. This is interpolation on six points, so it does not assert that R is a polynomial. Direct finite-difference calculation yields

$$\beta_5 = D_5/120, \quad \beta_3 = D_3/6 - D_5/48. \quad (\text{III.88.2})$$

In fact $D_3 = 6\beta_3 + 15\beta_5$ and $D_5 = 120\beta_5$.

Let $b_3 = -\beta_3$, $b_5 = -\beta_5$ be the target cubic and quintic coefficients that the core must supply. Its linear coefficient can always be set by a unique direction in the a coordinate. If the quintic direction is h , it must satisfy $N_5(h) = b_5$. When $h \neq 0$, its contribution to the cubic coefficient at quintic center v is

$$N_5(h)e_2(v/h) = b_5 e_2(v/h).$$

The remaining cubic direction must therefore lie in the norm fiber

$$N_3(h_3) = b_3 - b_5 e_2(v/h).$$

The cubic norm has one preimage of zero and L_3 preimages of every nonzero scalar; write this fiber size as

$$\nu_3(t) = 1 + (q^2 + q)1_{\{t \neq 0\}}.$$

If $b_5 = 0$, the quintic direction is zero, and the reflection fiber has size $\nu_3(b_3)$ at every core center. If $b_5 \neq 0$, there are L_5 possible quintic directions. For each such direction, $v \mapsto v/h$ is a bijection of E . Consequently the sum of the fiber size over all q^5 quintic centers is exactly

$$\begin{aligned} H(b_3, b_5) &= L_5 \sum_{x \in E} \nu_3(b_3 - b_5 e_2(x)) \\ &= L_5 \left[q^5 + (q^2 + q) \left(q^5 - q^4 - q^2 \chi(2b_3/b_5) \right) \right] \\ &= L_5 \left[q^7 - (q^4 + q^3) \chi(2b_3/b_5) \right]. \end{aligned}$$

Thus the mean reflection fiber over all core centers is

$$\mu(b_3, b_5) = \begin{cases} 1 + (q^2 + q)1_{\{b_3 \neq 0\}}, & b_5 = 0, \\ L_5 [q^2 - (q^{-1} + q^{-2})\chi(2b_3/b_5)], & b_5 \neq 0. \end{cases} \quad (\text{III.88.3})$$

This is an average; it need not be an integer. The integer quantity H above is the sum over quintic centers. It is not claimed that the fiber has size μ at each individual center.

Proof of Theorem III.88.2. The coefficients in the cubic and linear core centers affect only the unique final linear adjustment and do not change the preceding count. There are q^4 choices of these centers. Therefore the total number of symmetric center–direction pairs, including zero directions, is

$$q^9 \sum_{b,g} \mu(-\beta_3(b, g), -\beta_5(b, g)).$$

By (III.88.2), when $D_5 \neq 0$ the character in (III.88.3) is

$$\chi(2\beta_3/\beta_5) = \chi(40D_3/D_5 - 5).$$

Split the sum into $D_3 = D_5 = 0$, the T_0 cases, and the T_1 cases. Subtract the exactly q^{9+e} zero-direction pairs, one at each center. This gives (III.88.1).

For $q \geq 7$, every value of the second line of (III.88.3) is greater than one. The first line equals one precisely when $b_3 = 0$. Hence equality with the baseline holds precisely when $D_3 = D_5 = 0$ for every (b, g) , which is equivalent to vanishing of all pure third differences: a fifth difference is a second difference of a third difference.

For completeness, we prove the prime-field classification. Fix an $\mathbb{F}_p$ -basis of the additive domain $\mathbb{F}_q^e$, so it becomes $\mathbb{F}_p^{re}$. The function R has a unique reduced polynomial presentation with coefficients in $\mathbb{F}_q$,

$$R(X) = \sum_a c_a X^a, \quad 0 \leq a_i < p.$$

The restriction of R to each prime-field line has reduced degree at most two. To see this, a reduced one-variable polynomial of degree $d \in \{3, \dots, p-1\}$ has a nonzero third forward difference: its leading coefficient is multiplied by $d(d-1)(d-2) \neq 0$. Thus the assumed vanishing of every pure third difference excludes such a degree.

It follows that the coefficient of T^3 in $R(X + TH)$, after reducing T modulo $T^p - T$, vanishes as a function of (X, H) on $\mathbb{F}_p^{re} \times \mathbb{F}_p^{re}$. Suppose $c_a \neq 0$ and $|a| \geq 3$. Choose a multi-index $b \leq a$ with $|b| = 3$. This coefficient contains

$$c_a \binom{a}{b} X^{a-b} H^b.$$

Every X - and H -coordinate exponent is already below p ; moreover $\binom{a}{b} \neq 0$ in characteristic p . Different pairs (a, b) give distinct monomials $X^{a-b} H^b$. Contributions with larger $|b| \equiv 3 \pmod{p-1}$, which arise on reducing the parameter T , also have different H -monomials. The displayed term therefore cannot cancel. Uniqueness of the reduced polynomial as a function of (X, H) gives a contradiction. Hence R has total degree at most two in the prime-field coordinates.

Conversely, every such quadratic has zero pure third differences. Because p is odd, its quadratic part is $B(x, x)$ for a symmetric $\mathbb{F}_p$ -bilinear map B , and its degree-one part is an $\mathbb{F}_p$ -linear map L . This proves the classification, including proper extension fields and the zero-dimensional domain. $\square$

III.88.4 Explicit added monomials and exact finite checks

For one added coordinate $w \in \mathbb{F}_q$, all three elementary monomial extensions have closed formulas.

Corollary III.88.4 (Three exact monomial counts). *Write $P_j(a, u, v, w) = a + N_3(u) + N_5(v) + w^j$. Then*

$$\begin{aligned} A_6^{\text{sym}}(P_3) &= q^{10}(q-1)L_3, \\ A_6^{\text{sym}}(P_4) &= q^{12}(q-1), \\ A_6^{\text{sym}}(P_5) &= q^7(q-1)L_5[q^5 - (q^2-1)\chi(5)]. \end{aligned}$$

Proof. The statistics in Theorem III.88.2 are

$R(w)$	T_0	T_1	S
w^3	$q(q-1)$	0	0
w^4	$(q-1)^2$	0	0
w^5	0	$q(q-1)$	$(q-1)^2\chi(5)$.

For $R = w^3$, we have $D_3 = 6g^3$ and $D_5 = 0$; for $R = w^4$, we have $D_3 = 24bg^3$ and $D_5 = 0$. These give the first two rows. For $R = w^5$, the interpolated coefficients are $\beta_5 = g^5$ and $\beta_3 = 10b^2g^3$. Thus $D_5 \neq 0$ exactly when $g \neq 0$, and the character $\chi(2\beta_3/\beta_5) = \chi(20(b/g)^2)$ is zero for $b = 0$ and is $\chi(5)$ whenever $b, g \neq 0$. This proves the last row. Substitution into (III.88.1) and simplification gives the displayed counts. $\square$

The standard-library program `code/verify_six_ap_defect.py` constructs irreducible finite-field presentations and forms e_2 independently as the sum of pairwise Frobenius conjugates. It exhausts all $7^5 + 13^5 = 388100$ quintic-field elements, checks the entire e_2 distribution, and directly verifies cubic norm fibers on all $7^3 + 13^3 = 2540$ cubic-field elements. It then verifies all $7 \cdot 6 + 13 \cdot 12 = 198$ pairs of cubic targets and nonzero quintic targets by an independent weighted-fiber sum, including all three quadratic-character values. For example, for $q = 7$ the three integer sums H are

$$H_{+1} = 2299057999, \quad H_0 = 2306743943, \quad H_{-1} = 2314429887.$$

For $q = 13$ they are

$$H_{+1} = 1940550181219, \quad H_0 = 1941501864497, \quad H_{-1} = 1942453547775.$$

Finally it evaluates the added monomials at the actual six progression nodes and verifies the aggregate formula. The resulting ordered nontrivial progression counts are

$R(w)$	$q = 7$	$q = 13$
w^3	96606535158	302737248100404
w^4	83047723206	279577021469772
w^5	233281014955590	8654306689055041404.

These are exact factorized counts. The script does not enumerate all center–direction pairs in the ten-dimensional domain, and no such enumeration is needed for the proof.

What remains. For an eight-term core, the next norm layer has degree seven. The same descending procedure forces joint value distributions of several polarized coefficients, rather than one quadratic form. Determining whether those distributions admit comparably explicit character formulas is a concrete next problem. The global quintic avoidance dimension for symmetric four-term progressions, between four and eight in the preceding reports [173], remains unresolved here.

III.89 Source 45’s chapter: the real fourth-uniformity gap [write]

Added in the eighth write, 7 October 2026.

The next two sections are source 45’s Sections 4–5. *Logically they continue the chapter on the odd-order constants* (Section III.66), after source 34’s Section 10 and source 35’s Section 6; they are printed here because that chapter is no longer the last of Part III and no number printed earlier may move.

What is proved. For real centered f on a finite abelian group of odd order, $\|f\|_{U^4}^{16} \geq (4795/832) \|f\|_{U^2}^{16}$ (Theorem III.90.1), so $c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$, the best upper bound among sources 01–46 (source 49, placed by `f8bc5e2ec` and not yet written in, records $29/4$, that is $c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ (for its write)). The proof applies a mean-sensitive U^3 bound $Q_3(g) \geq v^2 + 10vw + 2w^2$ (Lemma III.90.4; the three-pairing count of sources 21, 35 and 38 with the zero frequency kept) to the derivatives $f \cdot f(\cdot + h)$, the even-moment bounds of the autocorrelation (sources 29’s and 34’s C_m), $W \geq V - 1$, and the scalar certificate $P(t) = -t^2/4 + (33/26)t - 99/676$ with $10t^2 - 4P = 11(t - 3/13)^2$. The five-point function $(2, 1, -2, -2, 1)$ on $\mathbb{Z}/5\mathbb{Z}$ gives $c_4^{\text{odd}} \geq (39037448/404304625)^{1/4} = 0.5574336$ (Proposition III.90.9), weaker than sources 28 and 34 (0.5587294) and the batch-121 certificate (0.5587725); a dual-cube identity rules out every real single-character function of order at least five, order three treated separately (sources 28’s and 34’s results, uncited). Section III.91 multiplies source 04’s quartic coefficient P_d by $(832/4795)^{1/4}$ for $d \geq 4$ (Corollary III.91.2) and computes the defect of rectified holes. The bracket is now $0.5587725 \leq c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$ (item III.1).

Relation to the formal statements (checked at HEAD 6d4942222). `lemma_3_10` (Sections01_03.lean, line 340) has the exact companion `lemma_3_10_holds` (Proofs03CubeUpper.lean, line 82); no declaration concerns c_d^{odd} . Nothing in this chapter is formalized.

III.90 A quantitative real fourth-uniformity gap

Source 45, delivered Section 4.

The real, centered setting retains information that is lost in the ordinary monotonicity of Gowers norms. The pinned report [202] establishes the sharp comparison

$$\|f\|_{U^3}^8 \geq 2 \|f\|_{U^2}^8$$

for real functions of mean zero on finite abelian groups of odd order. It also asks for the corresponding optimal comparison with U^d when $d \geq 4$. We give a strictly stronger fourth-order comparison, an exact finite lower bound for its optimal constant, and a structural obstruction to carrying the third-order extremizers into dimension four.

Our argument uses the Fourier and cube identities underlying Gowers’s proof [1], together with an elementary moment count and a rational inequality certificate. The broader theory of functions whose Gowers norms are near an extremal value under L^p or L^∞ constraints is developed, for example, by Eisner and Tao [7]. The quotient considered here imposes realness and centering and compares two different uniformity norms. We state the exact remaining uncertainty: the numerical constant proved below is a certified lower bound for the fourth-order norm gap, and its optimal value is not determined in this article.

III.90.1 Normalization and the results

Let G be a finite abelian group, written additively. Averages over G and its powers are normalized. Characters take values on the unit circle, and the dual group $\widehat{G}$ is also written additively. For a function $f : G \rightarrow \mathbb{C}$, write

$$a_\xi = \widehat{f}(\xi) = \mathbb{E}_{x \in G} f(x) \overline{\xi(x)}, \quad f(x) = \sum_{\xi \in \widehat{G}} a_\xi \xi(x). \quad (\text{III.90.1})$$

The Fourier sums in this section are *unnormalized*; in particular, the normalization is already incorporated in a_ξ .

Let $\mathcal{C}$ denote complex conjugation. For $d \geq 1$ set

$$Q_d(f) = \|f\|_{U^d}^{2d} = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{\omega \in \{0,1\}^d} \mathcal{C}^{|\omega|} f(x + \omega \cdot h). \quad (\text{III.90.2})$$

Here $\|f\|_{U^1} = |\mathbb{E}f|$; for $d \geq 2$ this is a norm. For real functions, every conjugation in (III.90.2) can be omitted. The recursion, with $\Delta_h f(x) = f(x) \overline{f(x+h)}$, is

$$Q_{d+1}(f) = \mathbb{E}_h Q_d(\Delta_h f). \quad (\text{III.90.3})$$

Orthogonality gives

$$S(f) := Q_2(f) = \sum_{\xi} |a_\xi|^4. \quad (\text{III.90.4})$$

If f is real, $a_{-\xi} = \overline{a_\xi}$; if it is centered, $a_0 = 0$. Odd group order ensures that every nonzero character belongs to a distinct pair $\{\xi, -\xi\}$.

Theorem III.90.1 (Explicit real fourth-order gap). *Let G be a finite abelian group of odd order and let $f : G \rightarrow \mathbb{R}$ satisfy $\mathbb{E}f = 0$. Then*

$$\boxed{Q_4(f) \geq \frac{4795}{832} S(f)^4}. \quad (\text{III.90.5})$$

Consequently, for every $d \geq 4$,

$$\|f\|_{U^2}^4 \leq \left(\frac{832}{4795} \right)^{1/4} \|f\|_{U^d}^4. \quad (\text{III.90.6})$$

No bound on $\|f\|_\infty$ is required.

[write] **7 October 2026; sources 40–46.** *Where this stands* (checked at the write). In source 29's normalization $\|f\|_{U^4}^{16} \geq C \|f\|_{U^2}^{16}$ the coefficients are $(323/160)^2$ (28), $35/8$ (29), 5 (35), $81/16$ (34) and $4795/832 = 5.7632212$ (here), that is $c_4^{\text{odd}} \leq 0.7038153, 0.6914416, 0.6687403, 2/3$ and $(832/4795)^{1/4} = 0.6454070$; $4795/832 - 81/16 = 583/832$. So this is the best upper bound among sources 01–46 (source 49, placed by **f8bc5e2ec** and not yet written in, records $29/4$, that is $c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ (for its write)), and $0.5587725 \leq c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$, the lower end being the batch-121 intake's certificate on $\mathbb{Z}/7\mathbb{Z}$ (note after source 28's Corollary III.70.8). The intake re-derived the proof and checked every identity of the certificate in exact rationals $(10t^2 - 4P = 11(t - 3/13)^2$, the final quadratic in V) and every intermediate inequality on 47 functions over nine odd groups. Local numerical minimization puts the truth near $C \approx 10.2571$ (dossier 122; evidence only).

[write] **7 October 2026; sources 47–53.** *Superseded by source 49* (now printed, Section III.92): its Theorem III.93.1 has the coefficient $29/4 > 4795/832$ (the difference is $1237/832$), so $c_4^{\text{odd}} \leq (4/29)^{1/4} =$

0.6094183 and the bracket among sources 01–53 is $0.5587725 \leq c_4^{\text{odd}} \leq 0.6094183$; source 49 was written before this source arrived and does not cite it. The statement “best among sources 01–46” above stays true for those sources.

The theorem is homogeneous. Thus it also applies to the 1-bounded functions used in density arguments, and imposing $\|f\|_\infty \leq 1$ does not alter the optimal quotient below: every nonzero function on a finite group can be rescaled to satisfy that condition.

Define

$$c_d^{\text{odd}} = \sup \left\{ \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4} : \begin{array}{l} G \text{ finite abelian of odd order,} \\ f : G \rightarrow \mathbb{R}, \quad \mathbb{E}f = 0, \quad f \neq 0 \end{array} \right\}. \quad (\text{III.90.7})$$

Corollary III.90.2 (A rigorous interval for the optimal constant). *One has*

$$\left(\frac{39037448}{404304625} \right)^{1/4} \leq c_4^{\text{odd}} \leq \left(\frac{832}{4795} \right)^{1/4}. \quad (\text{III.90.8})$$

Numerically, the two endpoints are

$$0.5574336443118301 \quad \text{and} \quad 0.6454070108505990.$$

Moreover, no function of the form $2\text{Re}(c\xi)$, with $c \neq 0$ and ξ a nontrivial character, attains the supremum defining c_4^{odd} .

The upper endpoint is proved in Section III.90.4. The exact five-point example providing the lower endpoint appears in Section III.90.6. The exclusion of the single-character class is proved in Section III.90.5.

Each of the three restrictions in Theorem III.90.1 matters. A nontrivial complex character is centered and has U^j norm equal to 1 for every $j \geq 2$, so realness cannot be dropped. A nonzero constant function shows that centering cannot be dropped. On $\mathbb{Z}/2\mathbb{Z}$, the nontrivial real character is centered and again has U^j norm equal to 1 for every $j \geq 2$, so an unrestricted even-order version is false.

For completeness, ordinary monotonicity follows directly from the cube identities. With $j - 1$ derivatives grouped into F_h , one has

$$Q_j(f) = \mathbb{E}_h |\mathbb{E}_x F_h(x)|^2, \quad Q_{j+1}(f) = \mathbb{E}_h \|F_h\|_{U^2}^4.$$

By (III.90.4), $\|F_h\|_{U^2}^4 \geq |\mathbb{E}F_h|^4$, and Cauchy–Schwarz therefore gives $Q_{j+1}(f) \geq Q_j(f)^2$. This proves $\|f\|_{U^{j+1}} \geq \|f\|_{U^j}$ and the passage from dimension four to (III.90.6). Combining only monotonicity with the third-order result in [202] would give the weaker constant $1/\sqrt{2}$.

III.90.2 A third-order inequality retaining the mean

The essential first step is to retain the zero Fourier coefficient when performing the third-order argument. Derivatives of a centered function need not be centered; their means are its autocorrelations. Discarding those means before the next averaging would lose the improvement.

Lemma III.90.3 (Fourier face identity). *For every complex function g on a finite abelian group, with normalized Fourier coefficients a_r , one has*

$$Q_3(g) = \sum_{s,t \in \widehat{G}} |T(s,t)|^2, \quad (\text{III.90.9})$$

$$T(s,t) = \sum_{r \in \widehat{G}} a_r \overline{a_{r+s}} \overline{a_{r+t}} a_{r+s+t}. \quad (\text{III.90.10})$$

The identity requires no assumption on the order of the group.

Proof. Expand (III.90.2) for $d = 3$ using Fourier inversion. For the term with frequency labels r_ω , orthogonality in x, h_1, h_2, h_3 imposes exactly

$$\sum_{\omega} (-1)^{|\omega|} r_\omega = 0, \quad \sum_{\omega} (-1)^{|\omega|} \omega_i r_\omega = 0 \quad (1 \leq i \leq 3). \quad (\text{III.90.11})$$

These equations are equations in $\widehat{G}$, so they must be solved without dividing by an integer. Subtract the h_3 equation from the x equation. The result is

$$r_{000} - r_{100} - r_{010} + r_{110} = 0.$$

Subtracting the h_2 and h_1 equations in the same way determines r_{101} and r_{011} . Substitution in any one of the h_i equations then determines r_{111} . Thus, writing

$$r = r_{000}, \quad s = r_{100} - r, \quad t = r_{010} - r, \quad u = r_{001} - r,$$

all solutions are exactly

$$r_\omega = r + \omega_1 s + \omega_2 t + \omega_3 u.$$

Conversely these labels satisfy (III.90.11), as direct substitution shows. There are no multiplicities or divisibility factors.

For fixed s, t , put

$$P_{s,t}(r) = a_r \overline{a_{r+s}} \overline{a_{r+t}} a_{r+s+t}.$$

The two parallel frequency faces contribute $P_{s,t}(r) \overline{P_{s,t}(r+u)}$. Summing over r, u is therefore $|\sum_r P_{s,t}(r)|^2$. The remaining sum over s, t proves the identity. All averages used in orthogonality equal either zero or one; this accounts for the unnormalized sums in (III.90.9). $\square$

Lemma III.90.4 (Mean-sensitive real third-order bound). *Let G have odd order and let $g : G \rightarrow \mathbb{R}$ be arbitrary. Set*

$$v = |\mathbb{E}g|^4, \quad w = Q_2(g) - v \geq 0. \quad (\text{III.90.12})$$

Then

$$\boxed{Q_3(g) \geq v^2 + 10vw + 2w^2.} \quad (\text{III.90.13})$$

Proof. Put

$$\lambda_r = |a_r|^2, \quad H = \sum_r \lambda_r^2 = v + w, \quad M = \sum_r \lambda_r^4.$$

Because g is real, $\lambda_{-r} = \lambda_r$ and $\lambda_0^2 = v$. Introduce the unnormalized additive energy of the nonnegative weights λ :

$$\mathcal{E}(\lambda) = \sum_{r+s=t+u} \lambda_r \lambda_s \lambda_t \lambda_u.$$

On either coordinate axis in (III.90.9),

$$T(0, t) = \sum_r \lambda_r \lambda_{r+t}.$$

The squared sums on one axis add up to $\mathcal{E}(\lambda)$. Retaining the union of the two axes, and subtracting their common origin once, gives

$$Q_3(g) \geq 2\mathcal{E}(\lambda) - H^2. \quad (\text{III.90.14})$$

Every summand of $\mathcal{E}(\lambda)$ is nonnegative. Retain the union of these three families of solutions of $r + s = t + u$:

$$\begin{aligned}\mathcal{A}: & \quad r = t, \quad s = u, \\ \mathcal{B}: & \quad r = u, \quad s = t, \\ \mathcal{D}: & \quad s = -r, \quad u = -t.\end{aligned}$$

Each family has total weight H^2 . Each pairwise intersection has total weight M . In the triple intersection all four indices coincide and $2r = 0$. Odd group order makes $r = 0$ its only possibility, so the triple intersection has weight $\lambda_0^4 = v^2$. Inclusion–exclusion yields

$$\mathcal{E}(\lambda) \geq 3H^2 - 3M + v^2. \quad (\text{III.90.15})$$

For representatives of the nonzero opposite pairs write $b_j = \lambda_{r_j}^2$. Then

$$w = 2 \sum_j b_j, \quad M = v^2 + 2 \sum_j b_j^2 \leq v^2 + \frac{1}{2}w^2.$$

Substitution into (III.90.14) and (III.90.15) gives

$$\begin{aligned}Q_3(g) &\geq 5(v+w)^2 - 6\left(v^2 + \frac{1}{2}w^2\right) + 2v^2 \\ &= v^2 + 10vw + 2w^2,\end{aligned}$$

as required. □

When g is centered this reduces to the third-order factor 2. The extra mixed term in (III.90.13) is what permits the fourth-order improvement. Notice also why the zero frequency could not be omitted in the proof: it contributes the triple intersection v^2 in (III.90.15).

III.90.3 Positive Fourier coefficients and even autocorrelation moments

For a real centered function f , define its autocorrelation by

$$R(h) = \mathbb{E}_x f(x)f(x+h). \quad (\text{III.90.16})$$

Fourier inversion and realness give

$$R(h) = \sum_{\xi} \lambda_{\xi} \xi(h), \quad \lambda_{\xi} = \left| \widehat{f}(\xi) \right|^2, \quad \lambda_{-\xi} = \lambda_{\xi}, \quad \lambda_0 = 0. \quad (\text{III.90.17})$$

Thus R is real, even, and has nonnegative Fourier coefficients. Also

$$\mathbb{E}_h R(h)^2 = \sum_{\xi} \lambda_{\xi}^2 = S(f). \quad (\text{III.90.18})$$

Lemma III.90.5 (All even moments of the autocorrelation). *If G has odd order, $f : G \rightarrow \mathbb{R}$ is centered, and $k \geq 1$ is an integer, then*

$$\boxed{\mathbb{E}_h R(h)^{2k} \geq \frac{\binom{2k}{k}}{2^k} S(f)^k.} \quad (\text{III.90.19})$$

For $k = 1$ equality always holds.

Proof. Choose one character from each nonzero opposite pair. With $b_j \geq 0$ we may write

$$R(h) = \sum_j b_j (\chi_j(h) + \chi_j(h)^{-1}), \quad S(f) = 2 \sum_j b_j^2.$$

In the expansion of $\mathbb{E}R^{2k}$, each character word averages to zero or one. Its scalar coefficient is nonnegative. Retain those words that contain the same number k_j of occurrences of χ_j and χ_j^{-1} , for every j . These words certainly average to one, irrespective of any other relations among the characters. Their total contribution is

$$(2k)! \sum_{\sum_j k_j = k} \prod_j \frac{b_j^{2k_j}}{(k_j!)^2}. \quad (\text{III.90.20})$$

For every multi-index in this sum, the multinomial coefficient $k!/\prod_j k_j!$ is a positive integer. Consequently

$$\prod_j \frac{1}{(k_j!)^2} \geq \frac{1}{k!} \prod_j \frac{1}{k_j!}.$$

Using the multinomial theorem in (III.90.20), we obtain

$$\begin{aligned} \mathbb{E}R^{2k} &\geq \frac{(2k)!}{k!} \sum_{\sum_j k_j = k} \prod_j \frac{b_j^{2k_j}}{k_j!} \\ &= \binom{2k}{k} \left(\sum_j b_j^2 \right)^k = \frac{\binom{2k}{k}}{2^k} S(f)^k. \end{aligned}$$

For $k = 1$, this is precisely (III.90.18). □

The lemma does not require the character representatives to be independent or dissociated. Additional additive relations contribute further nonnegative terms; none invalidates the selected family of balanced words. This observation is particularly useful for handling small odd torsion uniformly.

We shall use the four instances

$$\begin{aligned} \mathbb{E}R^2 &= S, & \mathbb{E}R^4 &\geq \frac{3}{2}S^2, \\ \mathbb{E}R^6 &\geq \frac{5}{2}S^3, & \mathbb{E}R^8 &\geq \frac{35}{8}S^4, \end{aligned} \quad (\text{III.90.21})$$

where $S = S(f)$. These particular numerical coefficients are not arbitrary estimates: they are $\binom{2k}{k}/2^k$ at $k = 1, 2, 3, 4$.

III.90.4 A rational certificate for the fourth-order gap

Proof of Theorem III.90.1. If $f = 0$, the result is immediate. Otherwise $S(f) > 0$ by Fourier inversion. Replace f by $S(f)^{-1/4}f$; it is enough to prove the statement under the normalization $S(f) = 1$.

For each $h \in G$ let

$$g_h(x) = f(x)f(x+h), \quad t(h) = R(h)^2, \quad w(h) = Q_2(g_h) - t(h)^2. \quad (\text{III.90.22})$$

The zero Fourier coefficient of g_h is $R(h)$, so $w(h) \geq 0$. Introduce the scalar moments

$$V = \mathbb{E}t^2, \quad B = \mathbb{E}t^3, \quad A = \mathbb{E}t^4, \quad W = \mathbb{E}w. \quad (\text{III.90.23})$$

By (III.90.21),

$$\mathbb{E}t = 1, \quad V \geq \frac{3}{2}, \quad B \geq \frac{5}{2}, \quad A \geq \frac{35}{8}. \quad (\text{III.90.24})$$

There is one more constraint linking W and V . The cube recursion gives

$$V + W = \mathbb{E}_h Q_2(g_h) = Q_3(f).$$

Apply the axis estimate (III.90.14) directly to f . Here its H equals $S(f) = 1$, and Fourier orthogonality with the even weights λ_ξ gives

$$\mathcal{E}(\lambda) = \mathbb{E}_h R(h)^4 = V.$$

Therefore

$$V + W \geq 2V - 1, \quad W \geq V - 1. \quad (\text{III.90.25})$$

Now apply Lemma III.90.4 to each g_h . Its variable v is $R(h)^4 = t(h)^2$, and its variable w is exactly that in (III.90.22). Averaging the result yields

$$Q_4(f) \geq \mathbb{E}[t^4 + 10t^2w + 2w^2]. \quad (\text{III.90.26})$$

We finish with a scalar inequality, retaining the mean and the next three moments of t . Set

$$a = \frac{99}{676}, \quad b = \frac{33}{26}, \quad P(t) = -\frac{1}{4}t^2 + bt - a. \quad (\text{III.90.27})$$

The exact identity

$$10t^2 - 4P(t) = 11 \left(t - \frac{3}{13} \right)^2 \quad (\text{III.90.28})$$

and $w \geq 0$ imply

$$t^4 + 10t^2w + 2w^2 \geq t^4 + 2(w + P(t))^2 - 2P(t)^2. \quad (\text{III.90.29})$$

The pointwise difference is $11w(t - 3/13)^2$, so the direction of this inequality does not depend on the sign of $P(t)$.

Expand the nonsquare terms:

$$t^4 - 2P(t)^2 = \frac{7}{8}t^4 + \frac{33}{26}t^3 - \frac{2277}{676}t^2 + 4abt - 2a^2. \quad (\text{III.90.30})$$

Taking expectations and applying Jensen only to the displayed square in (III.90.29), we obtain

$$\begin{aligned} Q_4(f) &\geq \frac{7}{8}A + \frac{33}{26}B - \frac{2277}{676}V + 4ab - 2a^2 \\ &\quad + 2 \left(W - \frac{1}{4}V + b - a \right)^2. \end{aligned} \quad (\text{III.90.31})$$

The coefficients of A and B are positive, permitting substitution of their lower bounds in (III.90.24). The expression in the last square is also positive, because

$$W - \frac{1}{4}V + b - a \geq \frac{3}{4}V + \frac{83}{676} > 0. \quad (\text{III.90.32})$$

It is therefore legitimate to replace W by its lower bound $V - 1$ in that square. After these substitutions, the right side of (III.90.31) is exactly

$$\frac{4795}{832} + \frac{3}{8} \left(V - \frac{3}{2} \right) + \frac{9}{8} \left(V - \frac{3}{2} \right)^2. \quad (\text{III.90.33})$$

This last equality is a rational polynomial identity, obtained by substituting the two fractions in (III.90.27); the supplementary verifier checks it using exact fraction arithmetic.

Since $V \geq 3/2$, both additional terms in (III.90.33) are nonnegative. We have proved $Q_4(f) \geq 4795/832$ under $S(f) = 1$. Restoring the scale proves (III.90.5), and monotonicity proves (III.90.6). $\square$

Several losses in this proof are visible and independently nonnegative: the discarded terms in the Fourier face sum, the unselected terms in the autocorrelation moments, $11w(t - 3/13)^2$, and the variance discarded by Jensen. Their equality conditions have not been jointly optimized. Accordingly, $4795/832$ is presented as an explicit improved certificate, not as the optimal fourth-order factor.

III.90.5 Why one Fourier pair ceases to be extremal

The exact third-order comparison in [202] is attained by real parts of characters. The fourth-order cube has more surviving frequency patterns, and these produce a first-order direction of improvement away from that family.

Let ξ have odd order $q \geq 3$ and fix $\theta \in \mathbb{R}$. For each integer j define the real function

$$\phi_j(x) = e^{ij\theta} \xi(x)^j + e^{-ij\theta} \xi(x)^{-j}. \quad (\text{III.90.34})$$

The index j is an integer exponent, not a claim that all these Fourier pairs are distinct. In particular, the possible aliases at orders five and seven will be treated explicitly by orthogonality.

For a real function u , define its fourth cube dual function by

$$\mathcal{D}_4 u(x) = \mathbb{E}_{h_1, h_2, h_3, h_4} \prod_{\substack{\omega \in \{0,1\}^4 \\ \omega \neq 0}} u(x + \omega \cdot h). \quad (\text{III.90.35})$$

Then $Q_4(u) = \mathbb{E}_x u(x) \mathcal{D}_4 u(x)$.

Lemma III.90.6 (Dual polynomial for a real character pair). *If $q \geq 5$ is odd, then*

$$\mathcal{D}_4 \phi_1 = 111\phi_1 + 27\phi_3 + \phi_5. \quad (\text{III.90.36})$$

In particular,

$$Q_2(\phi_1) = 2, \quad Q_4(\phi_1) = 222, \quad \mathbb{E} \phi_3 \mathcal{D}_4 \phi_1 = 54. \quad (\text{III.90.37})$$

Proof. Expand the fifteen factors in (III.90.35). A term is indexed by signs $\sigma_\omega \in \{-1, 1\}$, one for each nonzero vertex. If $j = \sum_{\omega \neq 0} \sigma_\omega$, that term contributes $e^{ij\theta} \xi(x)^j$ precisely when

$$\sum_{\omega \neq 0} \omega_i \sigma_\omega \equiv 0 \pmod{q} \quad (1 \leq i \leq 4). \quad (\text{III.90.38})$$

Each sum in (III.90.38) consists of eight signs and is an even integer between -8 and 8 . For odd $q \geq 5$, a nonzero multiple of q in this range cannot be even: the first possible nonzero even multiple has absolute value $2q \geq 10$. Thus the congruences are integer equalities, uniformly in all such q .

The number of integer-balanced sign assignments is recorded in Table III.90.1. An exact polynomial certificate for the count is

$$\begin{aligned} [z_1^4 z_2^4 z_3^4 z_4^4] \prod_{\omega \in \{0,1\}^4 \setminus \{0\}} \left(1 + y \prod_{i=1}^4 z_i^{\omega_i} \right) \\ = y^5 + 27y^6 + 111y^7 + 111y^8 + 27y^9 + y^{10}. \end{aligned} \quad (\text{III.90.39})$$

Indeed, a chosen subset of positive-sign vertices must contain four vertices with coordinate i equal to one, for each i . If the chosen subset has size k , its total sign sum is $j = 2k - 15$.

For clarity, the coefficient evaluation in (III.90.39) can be performed by the following finite integer recurrence. Begin with the coefficient 1 at $(k, u_1, u_2, u_3, u_4) = (0, 0, 0, 0, 0)$. For each successive nonzero vertex ω , copy each current coefficient to the unchanged index and add it at $(k + 1, u_1 + \omega_1, \dots, u_4 + \omega_4)$. Terms having some $u_i > 4$ may be discarded. After fifteen steps, the coefficients at $(k, 4, 4, 4, 4)$ are respectively 1, 27, 111, 111, 27, 1 for $k = 5, 6, 7, 8, 9, 10$, and zero otherwise. This specifies an exact integer evaluation of the displayed finite product; the supplied program also checks it independently by enumerating all 2^{15} sign assignments.

Total sign sum j	-5	-3	-1	1	3	5
Number of assignments	1	27	111	111	27	1

Table III.90.1: The fifteen-vertex integer-balanced sign count.

Grouping the opposite values of j proves (III.90.36), even when some resulting characters alias in the cyclic subgroup generated by ξ .

We now justify each orthogonality assertion required for (III.90.37). The two characters ξ and ξ^{-1} are distinct, so $Q_2(\phi_1) = 2$ and $\mathbb{E}\phi_1^2 = 2$. The exponents in $\phi_1\phi_3$ are $\pm 2, \pm 4$, and those in $\phi_1\phi_5$ are $\pm 4, \pm 6$. No odd $q \geq 5$ divides any of these nonzero exponents. Hence

$$\mathbb{E}\phi_1\phi_3 = \mathbb{E}\phi_1\phi_5 = 0, \quad Q_4(\phi_1) = 111 \cdot 2 = 222.$$

Similarly q does not divide 6, so $\mathbb{E}\phi_3^2 = 2$. The exponents in $\phi_3\phi_5$ are $\pm 2, \pm 8$; again no odd $q \geq 5$ divides them. It follows that $\mathbb{E}\phi_3\phi_5 = 0$ and

$$\mathbb{E}\phi_3\mathcal{D}_4\phi_1 = 27 \cdot 2 = 54.$$

In particular, at $q = 5$ the function ϕ_5 is constant, but its pairing with either centered ϕ_1 or ϕ_3 still vanishes. At $q = 7$ some exponent pairs are represented by negative residues; the divisibility checks just made show that this causes no additional nonzero pairing. No independence assumption about ϕ_1, ϕ_3, ϕ_5 has been used. $\square$

Proposition III.90.7 (A strict direction of improvement). *For every odd $q \geq 5$, the real centered functions $f_\varepsilon = \phi_1 - \varepsilon\phi_3$ satisfy*

$$Q_2(f_\varepsilon) = 2(1 + \varepsilon^4), \quad (\text{III.90.40})$$

$$Q_4(f_\varepsilon) = 222 - 864\varepsilon + O(\varepsilon^2) \quad (\varepsilon \rightarrow 0). \quad (\text{III.90.41})$$

Consequently $Q_2(f_\varepsilon)/Q_4(f_\varepsilon)^{1/4}$ is strictly larger than $2/222^{1/4}$ for all sufficiently small positive ε .

Proof. The character pairs $\{\xi, \xi^{-1}\}$ and $\{\xi^3, \xi^{-3}\}$ are disjoint: an intersection would imply $q \mid 2$ or $q \mid 4$. The second pair consists of two nontrivial, distinct characters because $q \nmid 3, 6$. Fourier inversion therefore gives (III.90.40) exactly, including when $q = 5$.

For any real functions u, v , differentiating the finite polynomial $Q_4(u + \varepsilon v)$ gives sixteen terms, according to the vertex at which v replaces u . Each equals $\mathbb{E}v\mathcal{D}_4u$. To see this directly, if that vertex is η , change variables by

$$x' = x + \eta \cdot h, \quad h'_i = (-1)^{\eta_i} h_i.$$

This is a bijection of G^5 and relabels the vertices by coordinatewise flipping at η . Since the functions are real, no conjugation changes the product. Thus

$$\left. \frac{d}{d\varepsilon} Q_4(u + \varepsilon v) \right|_{\varepsilon=0} = 16\mathbb{E}v\mathcal{D}_4u. \quad (\text{III.90.42})$$

Taking $u = \phi_1$, $v = -\phi_3$ and using Lemma III.90.6 gives the derivative $-16 \cdot 54 = -864$. This proves (III.90.41). The numerator in the quotient has derivative zero at the origin, while its positive denominator has a strictly negative derivative. The asserted improvement follows. $\square$

The notation $O(\varepsilon^2)$ here hides no limiting group argument: for each fixed function pair, $Q_4(f_\varepsilon)$ is an ordinary real polynomial of degree at most sixteen. The nonzero first derivative alone establishes a strict improvement on that same finite group.

Lemma III.90.8 (Order-three characters). *If ξ has order three, then*

$$Q_4(\phi_1) = 286 + 32 \cos(6\theta) + 16 \cos(12\theta) = 262 + 32 \left(\cos(6\theta) + \frac{1}{2} \right)^2. \quad (\text{III.90.43})$$

In particular $Q_4(\phi_1) \geq 262$, whereas $Q_2(\phi_1) = 2$.

Proof. At order three the integer-equality reduction used in (III.90.38) is unavailable, so expand the full sixteen-vertex cube instead. A sign assignment contributes precisely when its total sign sum and each of its four coordinate sign sums vanish modulo three. Its phase is $e^{ij\theta}$, where j is the total sign sum. Since j is even and $|j| \leq 16$, the only possible values are $-12, -6, 0, 6, 12$. Their counts are

$$\begin{array}{c|ccccc} j & -12 & -6 & 0 & 6 & 12 \\ \hline \text{count} & 8 & 16 & 286 & 16 & 8. \end{array} \quad (\text{III.90.44})$$

Here is an exact recurrence evaluating these five integers. Enumerate the sixteen vertices as $\omega^{(1)}, \dots, \omega^{(16)}$. Let $N_j(s, r)$ count sign choices on the first j vertices with total integer sign sum s and four coordinate sums $r \in (\mathbb{Z}/3\mathbb{Z})^4$. Begin with $N_0(0, 0) = 1$ and all other entries zero. At a new vertex ω , add each entry $N_j(s, r)$ to both

$$N_{j+1}(s+1, r+\omega) \quad \text{and} \quad N_{j+1}(s-1, r-\omega).$$

After sixteen steps, retain $N_{16}(s, 0)$ for $3 \mid s$; this gives exactly (III.90.44). The recurrence is the signed counterpart of (III.90.39) and uses only integer additions. Exhaustive enumeration of the 2^{16} choices in the supplementary program independently reproduces the same table.

The opposite values of j now combine to give the first expression in (III.90.43). Substituting $\cos(12\theta) = 2\cos(6\theta)^2 - 1$ gives its completed-square form. The second-order statement follows from the two distinct Fourier coefficients of modulus one. $\square$

Every nonzero real single-character function is a positive scalar multiple of some ϕ_1 in (III.90.34). Combining the last two lemmas shows

$$\sup_{\substack{G \text{ odd, } \xi \neq 0 \\ c \neq 0}} \frac{\|2 \operatorname{Re}(c\xi)\|_{U^2}^4}{\|2 \operatorname{Re}(c\xi)\|_{U^4}^4} = \frac{2}{222^{1/4}} = 0.5181335948243558 \dots \quad (\text{III.90.45})$$

The upper value is attained whenever the character order is at least five, but Proposition III.90.7 gives an improving perturbation. The order-three class has an even smaller upper bound $2/262^{1/4}$. Thus the same character-pair extremizer description that works at order three fails at order four.

III.90.6 An exact five-point certificate

We now give a fixed rational-valued function realizing a quotient larger than the whole class in (III.90.45). Its verification uses integer cube products and requires neither numerical Fourier transforms nor an optimization routine.

Proposition III.90.9 (A rational lower bound). *On $G = \mathbb{Z}/5\mathbb{Z}$, set*

$$(f(0), f(1), f(2), f(3), f(4)) = (2, 1, -2, -2, 1). \quad (\text{III.90.46})$$

Then $\mathbb{E}f = 0$ and

$$Q_2(f) = \frac{94}{25}, \quad Q_3(f) = \frac{19208}{625}, \quad Q_4(f) = \frac{6468874}{3125}. \quad (\text{III.90.47})$$

Consequently

$$\left(\frac{\|f\|_{U^2}^4}{\|f\|_{U^4}^4} \right)^4 = \frac{39037448}{404304625} > \frac{8}{111}. \quad (\text{III.90.48})$$

Proof. The five values sum to zero. For dimension d , enumerate the 5^{d+1} tuples $(x, h_1, \dots, h_d)$ and form the integer

$$p(x, h) = \prod_{\omega \in \{0,1\}^d} f(x + \omega \cdot h \bmod 5).$$

Let $H_d(z)$ be the number of tuples with $p(x, h) = z$. This defines a finite exact certificate:

$$\sum_z H_d(z) = 5^{d+1}, \quad Q_d(f) = 5^{-d-1} \sum_z z H_d(z). \quad (\text{III.90.49})$$

For $d = 2$, the nonzero entries of this histogram are

z	−8	−4	−2	1	4	8	16
$H_2(z)$	8	24	8	6	36	24	19.

Their weighted sum is 470, hence $Q_2 = 470/125 = 94/25$. For $d = 3$ the nonzero entries are

z	−32	−16	1	4	16	32	64	128	256
$H_3(z)$	96	80	8	32	168	64	128	16	33.

They have total multiplicity 625 and weighted sum 19208. The fourth-dimensional histogram is given in Table III.90.2. Its multiplicities sum to 3125, and their product-weighted sum is 6468874. Equation (III.90.49) therefore gives all three identities in (III.90.47).

Cube product z	Multiplicity $H_4(z)$
−2048	160
−1024	320
−512	160
−64	32
1	10
16	80
256	760
512	320
1024	800
2048	32
4096	320
16384	80
65536	51
Total multiplicity	3125

Table III.90.2: Exact fourth-cube certificate for (III.90.46).

For an explicit algorithm generating these tables, start each histogram empty; run $x, h_1, \dots, h_d$ independently through $0, 1, 2, 3, 4$; initialize a product to one; for every binary vector ω multiply it by the corresponding entry of (III.90.46); and increase the resulting product's histogram count by one. This enumerates every normalized cube tuple exactly once and performs only exact integer operations. The complete implementation is included with the article.

Finally, using (III.90.47) and reducing the fraction gives

$$\frac{(94/25)^4}{6468874/3125} = \frac{39037448}{404304625}.$$

Cross-multiplication shows that this exceeds $8/111$, which is the fourth power of (III.90.45). This proves (III.90.48). $\square$

[write] **7 October 2026; sources 40–46.** Recomputed at the write in exact rationals: $Q_2 = 94/25$, $Q_4 = 6468874/3125$, $Q_2^4/Q_4 = 39037448/404304625$, fourth root 0.5574336 . Sources 28 and 34 give 0.5587294 and the batch-121 certificate 0.5587725 .

The function has only two nontrivial Fourier pairs. To see its connection with Proposition III.90.7, choose $\xi(x) = e^{2\pi ix/5}$ and $\theta = 0$. Its coefficients at the two opposite pairs have values

$$a_1 = \frac{5 + 3\sqrt{5}}{10}, \quad a_2 = \frac{5 - 3\sqrt{5}}{10}.$$

Here $a_2 < 0$, and on $\mathbb{Z}/5\mathbb{Z}$ the real functions ϕ_2 and ϕ_3 coincide at this phase. Thus

$$f = a_1(\phi_1 - \varepsilon\phi_3), \quad \varepsilon = -\frac{a_2}{a_1} = \frac{7 - 3\sqrt{5}}{2} > 0.$$

This identifies the exact rational example as a concrete member of the improving two-pair family. Its cube values, however, were established above without relying on radicals or trigonometric rounding. Dividing f by two gives a real centered 1-bounded example with the same quotient.

Proof of Corollary III.90.2. Theorem III.90.1 bounds every quotient in (III.90.7) by $(832/4795)^{1/4}$. Proposition III.90.9 supplies the lower endpoint. Its strict inequality (III.90.48), together with the full single-pair envelope (III.90.45), excludes every real single-character function as a global maximizer. This proves all parts of the corollary. $\square$

III.90.7 Reproducibility and the remaining problem

The mathematical upper bound is proved for every finite odd-order abelian group by the preceding inequalities. Its validity does not depend on a search over groups or functions. The finite computations serve two specific purposes: evaluating the explicit cosine sign-count identities and verifying the exact five-point example. The rational scalar certificate is also checked independently.

From the extracted package directory, run

```
python3 code/verify_u4.py --json certificates/u4_certificate.json
```

The program uses the Python standard library. Integer sums and `fractions.Fraction` carry all mathematical checks; floating-point approximations are unnecessary. It refuses to run under `python -O` or `python -OO`, so its assertions cannot be silently disabled. The deterministic JSON report records the rational certificate, the cosine coefficient tables, and the complete cube histograms and weighted sums. Direct sign enumeration and coefficient dynamic programming independently verify (III.90.39).

The following questions remain specific to this norm comparison. First, determine the exact value of c_4^{odd} and whether the supremum over all finite groups is attained. Second, classify the frequency support and phases of any extremizing functions: the single-pair classification valid at order three has now been excluded at order four. Third, retain more of the nonnegative remainders in (III.90.9), (III.90.20), and (III.90.29) to improve the explicit upper endpoint. Fourth, determine dimension-dependent bounds for c_d^{odd} , rather than using the same fourth-order bound for every $d \geq 4$.

The results here resolve the sharpness question for the single-character candidate and tighten a quantitative question posed in the pinned report. They do not determine the optimal fourth-order constant, assert a global publication-priority claim for the comparison, or provide a new global bound for progression-free sets. The exact scope of the progress is the proved interval (III.90.8), the universal inequality (III.90.5), and the strict improving direction in Proposition III.90.7.

III.91 Applications to cube counting and structured holes

Source 45, delivered Section 5.

We make the relationship with Gowers's Lemma 3.10 explicit. The cancellation and exact quartic coefficient in the next proposition are credited inputs from the previous ProveIt reports, especially sources 03, 04, 08, and 09 [202]. A short proof is included to make the new numerical substitution self-contained.

Let G be a finite abelian group of odd order, let $d \geq 2$, and put $q = 2^d$. For a real function W , write

$$Q_d(W) = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{v \in \{0,1\}^d} W(x + v \cdot h) = \|W\|_{U^d}^q.$$

Suppose $W = \rho + f$, where $\rho \geq 0$, $f : G \rightarrow \mathbb{R}$, and $\mathbb{E}f = 0$. Set $u = \|f\|_{U^d}$.

Proposition III.91.1 (Credited exact fourth-order expansion). *With the preceding notation,*

$$Q_d(\rho + f) = \rho^q + P_d \rho^{q-4} \|f\|_{U^2}^4 + R_d, \quad (\text{III.91.1})$$

$$P_d = \frac{6^d - 2 \cdot 4^d + 2^d}{8}, \quad (\text{III.91.2})$$

$$|R_d| \leq \sum_{r=5}^q \binom{q}{r} \rho^{q-r} u^r. \quad (\text{III.91.3})$$

The sum in (III.91.3) is zero if its index set is empty.

Proof. For a subset S of the Boolean vertices, expand the product by choosing f at the vertices in S and ρ elsewhere. Any three distinct Boolean vertices give jointly uniform affine forms on every finite abelian group. To verify this, reflect coordinates to move one vertex to zero. The two other row vectors are distinct nonzero Boolean vectors. Their supports supply a 2×2 minor of determinant 1 or -1 : if neither support contains the other use one coordinate from each difference; if they are nested use one coordinate in the smaller support and one in the difference. The one- and two-vertex cases are immediate. Thus every centered term of degree one, two, or three vanishes.

Four distinct Boolean vertices that are affinely dependent over $\mathbb{R}$ form an ordinary parallelogram. Indeed their affine span is a plane, because any three are independent, and projection onto two suitable coordinates is injective on that plane. The four projected points must be the four Boolean corners of a square. The affine inverse preserves their parallelogram relation. Any three of the corresponding forms are jointly uniform, so the fourth-order correlation is exactly $\|f\|_{U^2}^4$.

If the four vertices are affinely independent, reflect the first to zero and select a nonsingular 3×3 Boolean minor B from the other rows. Its determinant has absolute value at most two. For completeness, border B by a row and column of ones and replace its entries by $1 - 2B_{ij}$. Subtracting the first row shows that the new 4×4 determinant has absolute value $8|\det B|$, whereas Hadamard's inequality bounds it by 16. Hence $\det B = \pm 1$ or ± 2 . These integers act invertibly on an odd-order group; the adjugate identity makes the selected coordinate map invertible. Consequently the four forms are jointly uniform and their centered correlation is zero.

There are six solutions to $a + b = c + e$ in $\{0, 1\}$, hence 6^d ordered solutions on the Boolean cube. The solutions with identical unordered pairs on both sides number $2 \cdot 4^d - 2^d$. Every remaining solution uses four distinct vertices. Every parallelogram support has eight orderings and a unique pairing into opposite vertices, proving (III.91.2). Finally, cube Cauchy–Schwarz bounds a term supported on r vertices by u^r , since the other factors are one. Summing over supports of size at least five gives (III.91.3). $\square$

Corollary III.91.2 (Improved leading coefficient for $d \geq 4$). *Put $\kappa = (832/4795)^{1/4}$. If $d \geq 4$, then*

$$0 \leq Q_d(\rho + f) - \rho^q \leq \kappa P_d \rho^{q-4} u^4 + \sum_{r=5}^q \binom{q}{r} \rho^{q-r} u^r. \quad (\text{III.91.4})$$

Proof. The lower bound follows from $\|\rho + f\|_{U^d} \geq |\mathbb{E}(\rho + f)| = \rho$. The fourth-uniformity inequality of Section III.90 and monotonicity imply

$$\|f\|_{U^2}^4 \leq \kappa \|f\|_{U^4}^4 \leq \kappa \|f\|_{U^d}^4.$$

Substitute this into (III.91.1) and use the remainder bound. $\square$

The first values are $P_2 = 1$, $P_3 = 12$, $P_4 = 100$, and $P_5 = 720$. In dimension four, (III.91.4) has leading coefficient

$$100\kappa = 64.5407010850\dots,$$

where using only $c_3^{\text{odd}} = 1/\sqrt{2}$ would give $100/\sqrt{2} = 70.7106781187\dots$. This is an improvement of about 8.73% in that coefficient. The cancellation of the first three orders and the value $P_4 = 100$ were already present in the report; the smaller multiplier is the present contribution. The best possible quartic coefficient for this uniformity comparison is $P_d c_d^{\text{odd}}$, which is still unknown when $d \geq 4$.

III.91.1 The indicator and uniformity-parameter formulation

For $A \subset G$, let $N = |G|$, $\rho = |A|/N$, and $f = \mathbf{1}_A - \rho$. The number of ordered additive d -cubes in A , including degenerate cubes, is $N^{d+1}Q_d(\mathbf{1}_A)$. If $\|f\|_{U^d}^q \leq \alpha$, then for $d \geq 4$,

$$\#\{\text{ordered additive } d\text{-cubes in } A\} \leq N^{d+1} \left[\rho^q + \kappa P_d \rho^{q-4} \alpha^{4/q} + \sum_{r=5}^q \binom{q}{r} \rho^{q-r} \alpha^{r/q} \right]. \quad (\text{III.91.5})$$

The displayed uniformity condition is α -uniformity of degree $d - 1$ in the normalization of Gowers's lemma and the inspected repository definitions. This is a finite-parameter inequality, with an explicit remainder. It does not merely assert a formal asymptotic expansion. At fixed $\rho > 0$, the leading excess has exponent $4/q$ in α . Gowers's Minkowski estimate uses $N^{d+1}(\rho + \alpha^{1/q})^q$ [1, Lemma 3.10]; the improved exponent from centering is credited to the preceding reports, and (III.91.5) tightens its leading constant.

III.91.2 How far this reaches in the original proof

The torsion-free results refine the extremal discussion immediately before Gowers's Theorem 7.2. If $E(A) > M_m - 8(m - 3)$ and $m \geq 5$, they identify A exactly as a progression or a progression with an endpoint-adjacent hole. This is substantially more precise information in that narrow regime than a general large-subset structural conclusion. The hypothesis is correspondingly much closer to the maximum. It supplies no improved exponent in the general Balog–Szemerédi reduction.

The coset and collision results concern the regime where an additive triple almost always completes, and distinguish exact puncture extremizers from other high-energy sets. They provide useful rigidity statements for the energy-conversion step. They do not identify a large linear or multilinear piece under an arbitrary fixed positive energy density. In particular, the present collision equality theorem is a domain theorem; additional label-compatibility hypotheses must be retained when a graph or a derivative label is part of an application.

To obtain a new global bound for Szemerédi's theorem one would have to show that a sharper local estimate changes a parameter that survives all the later restriction, recurrence, and density-increment losses. That propagation is not supplied here. The proven improvements are the explicit local inequalities and classifications stated above.

III.91.3 A bridge between the two energy endpoints

The ordered energy theorem also describes the first defect levels when the hole of a dense set is rectified. A *Freiman 2-isomorphism* is a bijection that preserves, in both directions, every relation $a + b = c + d$ among four elements of its domain.

Corollary III.91.3 (Rectified holes). *Let C be a finite subgroup coset, $K \subset C$, $|K| = u \geq 5$, and $A = C \setminus K \neq \emptyset$, with $|A| = m$. Suppose that K admits a Freiman 2-isomorphism ι to a subset of a torsion-free abelian group. Put $\delta = u/m$ and*

$$\varepsilon_0(m, u) = \delta(1 - \delta) + \frac{u^3 - M_u}{m^3} = \delta(1 - \delta) + \frac{\delta^3}{3} - \frac{\delta}{3m^2}.$$

Then the actual defect $\varepsilon = 1 - E(A)/m^3$ obeys the following alternatives:

- (i) *If $\iota(K)$ is an arithmetic progression, then $\varepsilon = \varepsilon_0(m, u)$.*
- (ii) *If $\iota(K)$ is a second-level model from Section II.184, then $\varepsilon = \varepsilon_0(m, u) + 4(u - 2)/m^3$.*
- (iii) *In all other cases, $\varepsilon \geq \varepsilon_0(m, u) + 8(u - 3)/m^3$.*

All three lower values are attained, for example, when C is cyclic of order $m + u > 2u + 2$, by placing the three model holes in an interval without additive wraparound.

Proof. The complement identity proved in Section II.185 is exact:

$$\varepsilon = \delta(1 - \delta) + \frac{u^3 - E(K)}{m^3}.$$

Energy is preserved by ι . Substitute the three alternatives of the torsion-free theorem. For attainment, the integer holes $\{0, \dots, u - 1\}$, $H_u = \{0, \dots, u - 2, u\}$, and $J_u = \{0, \dots, u\} \setminus \{2\}$ lie in $[0, u]$. Their pair sums lie in $[0, 2u]$, so reduction modulo $m + u > 2u + 2$ preserves all additive quadruples. Their complements give the stated values. $\square$

For example, a rectified hole of relative size δ has a forced additional defect asymptotic to $\delta^3/3$ beyond the coset envelope, whereas a subgroup-coset hole has no such additional defect. This gives a precise way to distinguish two possible geometries of the deleted part. No rectification is inferred from the energy hypothesis alone.

III.92 Source 49's chapter: two-harmonic dual tests and $c_4^{\text{odd}} \leq (4/29)^{1/4}$ [write]

Added in the ninth write, 7 October 2026.

The next section is source 49's Section 7. *Logically it continues the chapter on the odd-order constants* (Section III.66) and source 45's chapter (Section III.89); it is printed here because no number printed earlier may move.

What is proved. For real centered f on a finite abelian group of odd order, $\|f\|_{U^4}^{16} \geq (29/4) \|f\|_{U^2}^{16}$, $\|f\|_{U^5}^4 \geq (50/29) \|f\|_{U^2}^4$, and $\|f\|_{U^4}^{16} \geq (15/2) \|f\|_{U^2}^{16}$ if a largest Fourier coefficient has order other than five (Theorem III.93.1): so $c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$, the best upper bound among sources 01–53, and $c_5^{\text{odd}} \leq 29/50$. The proof tests f against $g_t = \phi_1 + t\phi_3$ ($t = -1/7; -7/50$ at order five), whose fourth dual function is computed exactly by a finite integer certificate, and covers diffuse spectra by source 34's balanced-moment profile (Lemma III.93.2). It does not know source 45's 4795/832, which it supersedes (note after the theorem). The bracket is now $0.5587725 \leq c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ (item III.1); the lower end is the batch-121 certificate.

Relation to the formal statements (checked at HEAD 06b30ca52). `lemma_3_10` (Sections01_03.lean, line 340) has the exact companion `lemma_3_10_holds` (Proofs03CubeUpper.lean, line 82); no declaration concerns c_d^{odd} . Nothing in this chapter is formalized.

III.93 Two-harmonic dual tests and stronger real uniformity gaps

Source 49, delivered Section 7.

The real, centered norm comparison is a particularly concrete place where an explicit dual function improves a local constant in Gowers's framework. The earlier project reports prove the sharp inequality $\|f\|_{U^3}^8 \geq 2\|f\|_{U^2}^8$ and successively obtain $\|f\|_{U^4}^{16} \geq 5\|f\|_{U^2}^{16}$ and $\|f\|_{U^2}^4 \leq (2/3)\|f\|_{U^4}^4$; see [214, 213]. The last report also shows that a single Fourier pair is not extremal. Its negative third-harmonic example suggests a better *test function*, which is what we use here.

The new ingredient is the complete fourth dual function of $\phi_1 - \phi_3/7$, together with its small-order variants. Its unwanted Fourier coefficients are much smaller, relative to its principal coefficient, than those of the dual function of ϕ_1 alone. This permits a stronger comparison for an arbitrary function, not only for a two-harmonic example.

Theorem III.93.1 (Fourth and fifth real norm gaps). *Let G be a finite abelian group of odd order and let $f : G \rightarrow \mathbb{R}$ have mean zero. Then*

$$\|f\|_{U^4}^{16} \geq \frac{29}{4}\|f\|_{U^2}^{16}, \quad (\text{III.93.1})$$

$$\|f\|_{U^5}^4 \geq \frac{50}{29}\|f\|_{U^2}^4. \quad (\text{III.93.2})$$

If a Fourier coefficient of largest nonzero magnitude can be chosen at a character whose order is not five, then

$$\|f\|_{U^4}^{16} \geq \frac{15}{2}\|f\|_{U^2}^{16}. \quad (\text{III.93.3})$$

In particular, (III.93.3) holds whenever $5 \nmid |G|$. All three displayed inequalities are strict for nonzero f .

Consequently, with the earlier project's definition

$$c_d^{\text{odd}} = \sup_{G,f} \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4},$$

where G has odd order and f is nonzero, real, and centered, we have

$$c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183098\dots, \quad c_5^{\text{odd}} \leq \frac{29}{50} = 0.58. \quad (\text{III.93.4})$$

For groups of order prime to five the fourth constant is at most $(2/15)^{1/4} = 0.6042750794\dots$. These constants are not claimed to be sharp. In particular, the exact fourth-order extremal problem remains open. The fifth-order gain comes from its own autocorrelation moment combined with the new dual test, rather than merely propagating (III.93.1) by monotonicity.

[write] **7 October 2026; sources 47–53.** *Where this stands* (checked at the write). In source 29's normalization $\|f\|_{U^4}^{16} \geq C\|f\|_{U^2}^{16}$ the coefficients are $(323/160)^2$ (28), $35/8$ (29), 5 (35), $81/16$ (34), $4795/832$ (45) and $29/4$ (here; $15/2$ when a largest coefficient has order other than five), that is $c_4^{\text{odd}} \leq 0.7038153, 0.6914416, 0.6687403, 2/3, 0.6454070$ and $(4/29)^{1/4} = 0.6094183$ (respectively $(2/15)^{1/4} = 0.6042751$); $29/4 - 4795/832 = 1237/832$. So this is the best upper bound among sources 01–53, and $0.5587725 \leq c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$, the lower end being the batch-121 certificate (note after source 28's Corollary III.70.8). For $d = 5$, $c_5^{\text{odd}} \leq 29/50 = 0.58$ improves source 34's coupled bound $\beta_5 = 0.6099124$ (Remark III.71.10) and the monotone bound $c_5^{\text{odd}} \leq c_4^{\text{odd}} \leq 0.6094183$; since c_d^{odd} decreases in d , it also bounds c_d^{odd} for $d \geq 6$, where source 34's $\beta_6 = 0.5649044$, β_7 , β_8 and source 29's B_d for $d \geq 9$ are smaller. The intake recomputed every dual coefficient $d_{k,n}(t)$ ($n = 5, 7, 9, 11, \infty$) by an independent enumeration, and all seventeen rational comparisons exactly; its local numerical minima of the ratio stay above $29/4$ (near 10.2571, evidence only).

Proof status. The arguments below apply to every finite abelian group of odd order. One finite coefficient lemma is established by an exhaustive integer certificate: five expansions with $4^8 + 4^7 = 81\,920$ face assignments each. Every comparison is in rational arithmetic. Thus the coefficient lemma is a computer-assisted finite proof, not a numerical optimization or a sampling assertion. The accompanying `verify_dual_gap.py` regenerates the coefficients directly from the displayed finite formula, and `dual_gap_certificate.json` records them and the exact rational margins. No floating-point computation is used in the certificate. The decimals in the text are descriptive only.

III.93.1 Balanced moments and their dependence on a dominant pair

We use normalized counting measure and normalized Fourier coefficients $a_\xi = \mathbb{E}_x f(x) \overline{\xi(x)}$; Fourier sums are unnormalized. Write the dual group additively. For real centered f , $a_{-\xi} = \overline{a_\xi}$ and $a_0 = 0$. Nonzero frequencies occur in distinct inverse pairs because $|G|$ is odd. Put

$$S = \sum_{\xi} |a_\xi|^4 = \|f\|_{U^2}^4, \quad p_j = \frac{2|a_{\xi_j}|^4}{S}, \quad p = \max_j p_j,$$

when $f \neq 0$, choosing one representative from each nonzero inverse pair. Thus $\sum_j p_j = 1$. The selected largest pair will be denoted $\{\chi, -\chi\}$, and

$$a_\chi = Ae^{i\theta}, \quad A = (pS/2)^{1/4}.$$

The next lemma is the balanced-moment profile of [213]; we include its proof because it is an essential input to both new constants.

Lemma III.93.2 (Balanced-moment profile). *For an integer $m \geq 1$, define*

$$C_m = 2^{-m} \binom{2m}{m}, \quad H_m(p) = \sum_{k=0}^m \binom{m}{k}^2 p^k (1-p)^{m-k}.$$

For any selected Fourier pair of mass p ,

$$\mathbb{E}_h c_f(h)^{2m} \geq C_m S^m H_m(p), \quad c_f(h) = \mathbb{E}_x f(x) f(x-h). \quad (\text{III.93.5})$$

When $m = 2^{d-2}$ and $d \geq 2$, this implies

$$\frac{\|f\|_{U^d}^4}{S} \geq \{C_m H_m(p)\}^{1/m}. \quad (\text{III.93.6})$$

If the selected pair is largest and $p \leq 1/2$, then also

$$\left(\frac{\|f\|_{U^d}^4}{S} \right)^m \geq C_m \{m - (m-1)2^{-(m-1)}\}. \quad (\text{III.93.7})$$

For $p \in [1/2, 1]$, the right side of (III.93.6) is nonincreasing in p .

Proof. Set $b_j = |a_{\xi_j}|^2$. Fourier inversion gives $c_f(h) = \sum_{\xi} |a_\xi|^2 \xi(h)$. All monomials in its Fourier expansion have nonnegative coefficients. Retain those in which each ξ_j and $-\xi_j$ occurs the same number k_j of times. Their total contribution is

$$(2m)! \sum_{\sum k_j = m} \prod_j \frac{b_j^{2k_j}}{(k_j!)^2} = C_m S^m \sum_{\sum k_j = m} \binom{m}{k_1, k_2, \dots}^2 \prod_j p_j^{k_j}.$$

Group by the multiplicity k of the selected pair. In the remaining sum a squared multinomial coefficient is at least that coefficient. The multinomial theorem bounds this remaining sum below by $(1-p)^{m-k}$, proving (III.93.5).

The recursive identity and the elementary bound $\|g\|_{U^{d-1}} \geq |\mathbb{E}g|$ give

$$\|f\|_{U^d}^{2^d} = \mathbb{E}_h \|f(\cdot)f(\cdot-h)\|_{U^{d-1}}^{2^{d-1}} \geq \mathbb{E}_h |c_f(h)|^{2^{d-1}}.$$

This proves (III.93.6). For (III.93.7), interpret the normalized balanced sum as the expected multinomial coefficient of an m -trial multinomial distribution with probabilities p_j . That coefficient is one for a pure composition and at least m for every mixed composition. Hence the sum is at least $m - (m-1)\sum_j p_j^m$. Since $\sum_j p_j^m \leq p^{m-1}\sum_j p_j \leq 2^{-(m-1)}$, (III.93.7) follows.

Finally, coefficient extraction gives the identity

$$H_m(p) = \sum_{j=0}^{\lfloor m/2 \rfloor} \binom{m}{2j} \binom{2j}{j} \{p(1-p)\}^j.$$

Its coefficients are nonnegative, and $p(1-p)$ decreases on $[1/2, 1]$, proving the last assertion. $\square$

For the fourth norm, $m = 4$, and the profile specializes to

$$F(p) := \frac{35}{8} \{1 + 12p(1-p) + 6p^2(1-p)^2\} \leq \frac{\|f\|_{U^4}^{16}}{S^4}. \quad (\text{III.93.8})$$

If $p \leq 1/2$, (III.93.7) also gives the uniform lower bound $1015/64$ for the same ratio.

III.93.2 An exact dual polynomial and its modular aliases

For a real function g , let

$$D_4g(x) = \mathbb{E}_{h_1, h_2, h_3, h_4} \prod_{\substack{\omega \in \{0,1\}^4 \\ \omega \neq 0}} g(x + \omega \cdot h).$$

Gowers–Cauchy–Schwarz gives

$$|\mathbb{E}fD_4g| \leq \|f\|_{U^4} \|g\|_{U^4}^{15}, \quad \mathbb{E}gD_4g = \|g\|_{U^4}^{16}. \quad (\text{III.93.9})$$

These standard multilinear facts are part of the uniformity-norm machinery of [1].

For a character χ of odd order $n \geq 5$, write

$$\phi_j(x) = e^{ij\theta} \chi(x)^j + e^{-ij\theta} \chi(x)^{-j}, \quad g_t = \phi_1 + t\phi_3.$$

To specify the complete dual polynomial, introduce the Laurent polynomial

$$\mathcal{P}_t(X, Z) = \prod_{\substack{\omega \in \{0,1\}^4 \\ \omega \neq 0}} \left(XZ^\omega + X^{-1}Z^{-\omega} + tX^3Z^{3\omega} + tX^{-3}Z^{-3\omega} \right), \quad (\text{III.93.10})$$

where $Z^\omega = \prod_{j=1}^4 Z_j^{\omega_j}$. Let CT_n retain every monomial whose four Z exponents are divisible by n , and then set these Z monomials equal to one. Define

$$\text{CT}_n \mathcal{P}_t(X, Z) = \sum_k d_{k,n}(t) X^k.$$

The notation $n = \infty$ means the ordinary constant term in all four Z variables. Equivalently,

$$d_{k,n}(t) = \sum_{\substack{r_\omega \in \{-3,-1,1,3\} \ (\omega \neq 0) \\ \sum_\omega r_\omega = k \\ \sum_\omega \omega_j r_\omega \equiv 0 \pmod{n} \ (1 \leq j \leq 4)}} t^{\#\{\omega: |r_\omega|=3\}}. \quad (\text{III.93.11})$$

At $n = \infty$ the congruences mean integer equalities. Thus $d_{k,n}$ is an explicit integer polynomial of degree at most fifteen, and $d_{-k,n} = d_{k,n}$. Character orthogonality gives the exact identity

$$D_4 g_t = \sum_{k>0} d_{k,n}(t) \phi_k. \quad (\text{III.93.12})$$

Only positive odd k occur: fifteen odd integers have odd sum.

Lemma III.93.3 (The integer range and the small-order certificate). *For every odd $n \geq 13$, $d_{k,n} = d_{k,\infty}$ and these polynomials vanish for $|k| > 15$. For the choices of t in Table III.93.1, define M, E, W as follows. First, put*

$$M = d_{1,n}(t) - \sum_{\substack{k>0, \ k \neq 1 \\ k \equiv \pm 1 \pmod{n}}} |d_{k,n}(t)|, \quad (\text{III.93.13})$$

$$E = \sum_{\substack{k>0 \\ k \not\equiv 0,1,-1 \pmod{n}}} |d_{k,n}(t)|. \quad (\text{III.93.14})$$

To define W , collect the following exact expression as $b_0 + \sum_{\ell>0} b_\ell \cos(\ell\theta)$:

$$\begin{aligned} \|g_t\|_{U^4}^{16} = 2 \sum_{k>0} d_{k,n}(t) (&\mathbf{1}_{n|k-1} \cos((k-1)\theta) + \mathbf{1}_{n|k+1} \cos((k+1)\theta) \\ &+ t \mathbf{1}_{n|k-3} \cos((k-3)\theta) + t \mathbf{1}_{n|k+3} \cos((k+3)\theta)), \end{aligned} \quad (\text{III.93.15})$$

and put $W = b_0 + \sum_{\ell>0} |b_\ell|$. In the integer case, congruence means equality; thus $M = d_{1,\infty}(t)$, $E = \sum_{k>1} |d_{k,\infty}(t)|$, and $W = 2(d_{1,\infty}(t) + t d_{3,\infty}(t))$.

The table gives strict rational bounds $M > m_-/1000$, $E < e_+/1000$, and $W < w_+/1000$. Consequently

$$\begin{aligned} \frac{(2M)^{16}}{16W^{15}} &> \frac{242}{25}, \quad \frac{E}{M} < \frac{1}{200}, \quad n = 5, \\ \frac{(2M)^{16}}{16W^{15}} &> \frac{1033}{100}, \quad \frac{E}{M} < \frac{63}{10000}, \quad n = 7, 9, 11, \infty. \end{aligned} \quad (\text{III.93.16})$$

Proof. Each step exponent in (III.93.10) is a sum of eight odd integers of absolute value at most three. It is even and lies in $[-24, 24]$. For odd $n \geq 13$, divisibility by n therefore forces it to be zero as an integer. This proves agreement with the integer constant term. If all four step sums vanish, then

$$k = \sum_{\omega \neq 0} \left(1 - \frac{|\omega|}{2}\right) r_\omega.$$

The absolute values of the displayed coefficients sum to $4/2 + 4/2 + 1 = 5$, by grouping the vertices of weights one, two, three and four. Hence $|k| \leq 15$.

Formula (III.93.15) follows by inserting (III.93.12) into $\mathbb{E} g_t D_4 g_t$ and using character orthogonality. It gives the phase-uniform bound $\|g_t\|_{U^4}^{16} \leq W$.

n	t	m_-	e_+	w_+
5	$-7/50$	82581	389	165863
7	$-1/7$	82256	503	164459
9	$-1/7$	82255	509	164452
11	$-1/7$	82255	515	164452
∞	$-1/7$	82255	516	164452

Table III.93.1: Exact rational bounds for the two-harmonic dual certificate. The integer row applies to all odd character orders at least thirteen; its tail bound allows aliases and constants by overcounting them.

Here is the complete finite procedure establishing the table. Split the punctured four-cube into the upper face $(v, 1)$, $v \in \{0, 1\}^3$, and the seven nonzero lower-face vertices $(v, 0)$, $v \neq 0$. Enumerate the 4^8 upper assignments. Discard those whose total frequency is not zero modulo n , and group the rest by their three coordinate sums modulo n , their integer total frequency, and their number of frequencies of absolute value three. Enumerate the 4^7 lower assignments and match the negatives of their three coordinate sums to the upper groups. Add their integer totals and their counts of absolute-three frequencies. This counts each assignment in (III.93.11) exactly once: the fourth-step equation is the upper-face total, and the other three equations are exactly the matching conditions. For $n = \infty$, use integer sums throughout. There are no random choices and no omitted configurations.

The supplied verifier implements these two enumerations with integer counters, substitutes the rational value of t , forms (III.93.13)–(III.93.15), and verifies all table entries by integer comparisons after clearing denominators. It then verifies (III.93.16) using just the three rational bounds printed in each row. The JSON certificate includes every coefficient of every $d_{k,n}$, rather than only the successful comparisons. For illustration, the generic principal polynomial is

$$d_{1,\infty}(t) = 111 + 405t + 1785t^2 + 3705t^3 + 9615t^4 + 10934t^5 + 17225t^6 \\ + 13860t^7 + 13090t^8 + 6195t^9 + 4530t^{10} + 665t^{11} + 740t^{12},$$

and the exact generic norm check is

$$W|_{t=-1/7} = \frac{5465198316262956}{7^{16}}.$$

The same norm value appears in the earlier two-harmonic example [213]; the complete dual and the phase-uniform small-order bounds are the additional certificate used here. $\square$

The role of the order cutoff deserves emphasis. Although harmonics such as 11χ and 15χ may coincide up to sign at order thirteen, this does not invalidate the certificate. For odd $n \geq 13$ none of the odd frequencies $3, 5, \dots, 15$ is congruent to 1 or -1 modulo n : the relevant difference or sum is even and has absolute value at most sixteen, whereas the first nonzero even multiple of n is at least twenty-six. Tail–tail aliases are allowed in the sum of absolute values E , and constant harmonics have zero pairing with centered f . The integer row is thus a valid uniform bound for all these orders, including thirteen and fifteen.

III.93.3 A phase-uniform comparison for every Fourier profile

Proposition III.93.4 (Two-harmonic dual profiles). *Let $f \neq 0$ be real and centered on an odd-order group, and choose a largest pair of mass p and character order n . If $n \geq 5$, put $C = (2M)^{16}/(16W^{15})$*

and $\eta = E/M$ using Lemma III.93.3. Then

$$\frac{\|f\|_{U^4}^{16}}{S^4} \geq C[p^{1/4} - \eta(1-p)^{1/4}]_+^{16}. \quad (\text{III.93.17})$$

For $n \geq 13$ the integer row may be used. For all odd character orders, including three, the following common profile is valid:

$$\frac{\|f\|_{U^4}^4}{S} \geq \left(\frac{242}{25}\right)^{1/4} \left[p^{1/4} - \frac{63}{10000}(1-p)^{1/4}\right]_+^4. \quad (\text{III.93.18})$$

Proof. Choose the phase in the test function so that $a_\chi = Ae^{i\theta}$. Then $\mathbb{E}f\phi_1 = 2A$. Every other harmonic in the principal inverse pair has pairing of absolute value at most $2A$. For a nonconstant harmonic outside that pair,

$$|a_{k\chi}| \leq \{S(1-p)/2\}^{1/4}.$$

This bound holds separately for every occurrence, so no assumption of distinct tail frequencies is needed. Equations (III.93.12) and (III.93.13)–(III.93.14) give

$$\mathbb{E}fD_4g_t \geq 2M(S/2)^{1/4}\{p^{1/4} - (E/M)(1-p)^{1/4}\}.$$

Apply (III.93.9), use $\|g_t\|_{U^4}^{16} \leq W$, and raise to the sixteenth power when the bracket is positive. This proves (III.93.17); a nonpositive bracket gives the trivial lower bound zero. The factor sixteen in the definition of C comes from $(S/2)^4 = S^4/16$.

For an order-three character, average f over its kernel. Averaging translations is contractive for U^4 , since this is a translation-invariant norm. The resulting centered function on the three-point quotient is exactly $A\phi_1$. Direct character expansion of the full four-cube gives

$$\begin{aligned} \|\phi_1\|_{U^4}^{16} &= 286 + 32\cos(6\theta) + 16\cos(12\theta) \\ &= 262 + 32\{\cos(6\theta) + 1/2\}^2 \geq 262. \end{aligned}$$

For completeness, the counts at total frequencies $-12, -6, 0, 6, 12$ are $8, 16, 286, 16, 8$; the same exact verifier regenerates them from the 2^{16} sign assignments. Hence

$$\frac{\|f\|_{U^4}^{16}}{S^4} \geq \frac{131}{8}p^4 \quad (n=3). \quad (\text{III.93.19})$$

This order-three calculation was already established in [214, 213]. It dominates (III.93.18), because $131/8 > 242/25$ and subtracting the nonnegative tail term only decreases the latter profile. For $n \geq 5$, (III.93.18) follows directly from the bounds (III.93.16). This completes all possible character orders on an odd-order group. $\square$

There is also a convenient general formulation. For $d \geq 4$ and $m = 2^{d-2}$, the same proofs give the pointwise envelope

$$\frac{\|f\|_{U^d}^4}{S} \geq \max \left\{ \{C_m H_m(p)\}^{1/m}, \quad \left(\frac{242}{25}\right)^{1/4} \left[p^{1/4} - \frac{63}{10000}(1-p)^{1/4}\right]_+^4 \right\}. \quad (\text{III.93.20})$$

For $p \leq 1/2$, one may also take the maximum with the bound in (III.93.7). Order-specific constants from (III.93.16) improve the second entry. Formula (III.93.20) keeps the entire Fourier concentration parameter; it is stronger information than a single uniform constant.

III.93.4 Rational certificates for the uniform constants

Proof of Theorem III.93.1. The zero function is immediate. Assume $f \neq 0$, and write $R = \|f\|_{U^4}^{16}/S^4$. We first prove the stronger (III.93.3) when a selected largest pair has order different from five.

If $p \leq 1/2$, the diffuse bound gives $R \geq 1015/64 > 15/2$. If $1/2 \leq p \leq 15/16$, monotonicity of F gives

$$R \geq F(15/16) = \frac{1976905}{262144} > \frac{15}{2}.$$

Suppose $p \geq 15/16$. For a character of order three, (III.93.19) gives $R \geq (131/8)(15/16)^4 > 15/2$. Otherwise the character order is at least seven, and the appropriate row of (III.93.16) applies. The exact comparisons

$$(9839/10000)^4 < 15/16, \quad (1-p)^{1/4} \leq 1/2$$

show that the bracket in (III.93.17) is greater than

$$\frac{9839}{10000} - \frac{63}{10000} \cdot \frac{1}{2} = \frac{3923}{4000}.$$

Consequently

$$R > \frac{1033}{100} \left(\frac{3923}{4000} \right)^{16} > \frac{15}{2}.$$

The last rational lower bound is approximately 7.56891. This proves (III.93.3) without a restriction on the other Fourier coefficients.

It remains to treat a largest pair of order five. If $p \leq 1/2$, the preceding diffuse bound suffices. If $1/2 \leq p \leq 16/17$,

$$R \geq F(16/17) > \frac{29}{4}.$$

For $p \geq 16/17$, the order-five row gives $C > 242/25$ and $\eta < 1/200$. Since

$$(9849/10000)^4 < 16/17, \quad 1/17 < (1/2)^4,$$

the bracket in (III.93.17) is greater than

$$\frac{9849}{10000} - \frac{1}{200} \cdot \frac{1}{2} = \frac{614}{625}.$$

Thus

$$R > \frac{242}{25} \left(\frac{614}{625} \right)^{16} > \frac{29}{4}.$$

The last lower bound is approximately 7.28600. All these inequalities are included as exact fraction comparisons in the verifier. This proves the universal fourth-order statement. If $5 \nmid |G|$, no character can have order five, proving the stated corollary as well.

For the fifth norm put $R_5 = \|f\|_{U^5}^4/S$. Use (III.93.20) with $m = 8$. If $p \leq 1/2$, (III.93.7) is already greater than $50/29$ after taking the eighth root. If $1/2 \leq p \leq 99/100$, the balanced profile gives

$$R_5^8 \geq C_8 H_8(99/100) > (50/29)^8.$$

The left side of the last comparison is approximately 80.24185; the strict comparison itself is rational. Finally suppose $p \geq 99/100$. By monotonicity, $\|f\|_{U^5} \geq \|f\|_{U^4}$, and (III.93.18) applies. The rational inequalities

$$(997/1000)^4 < 99/100, \quad (8/25)^4 > 1/100$$

bound its bracket below by

$$\frac{997}{1000} - \frac{63}{10000} \cdot \frac{8}{25} = \frac{124373}{125000}.$$

Therefore

$$R_5^4 > \frac{242}{25} \left(\frac{124373}{125000} \right)^{16} > (50/29)^4.$$

This proves (III.93.2) and the asserted strictness. $\square$

What has and has not been settled. The new comparisons improve the project's explicit fourth- and fifth-order norm bounds, and the dominant-order condition identifies where the loss in the universal fourth-order certificate occurs. It does not prove that five-torsion worsens the true extremal constant. The lower constants supplied by actual two-harmonic examples remain below these upper bounds. An exact value of c_4^{odd} , a classification of its extremizers, and a uniform family of optimized dual tests at all orders remain open. The certificate deliberately states only the constants its exact rational comparisons establish.

III.93.5 How the two estimates cover the critical Fourier profiles

Figure III.93.1 displays the two competing lower estimates near the largest-pair masses used in the proof. The moment estimate falls as a single pair dominates, while the dual estimate rises. The proof uses rational thresholds on the two sides of this transition; the figure is explanatory and is not used to certify an inequality. The lower-bound constants and tail ratios in the plot are the printed rational constants, not numerical fits to coefficient data.




Figure III.93.1: For each applicable character order, the ratio is bounded below by the larger of its moment and dual profiles. The two estimates cover the transition region with a strict rational margin. Order three has its separate stronger projection bound.

III.94 Sources 63 and 60: the limit of the odd-order constants [write]

Added in the eleventh write, 7 October 2026.

The next five sections are source 63's Sections 2–5 and source 60's Section 2, on $c_d^{\text{odd}} = \sup \|f\|_{U^2}^4 / \|f\|_{U^d}^4$ over real centered f on groups of odd order, the constant of Section III.66. *Logically they continue that chapter* (sources 29, 28, 34, 35) and its continuations by sources 45 and 49 (Sections III.89 and III.92); they are printed here because no number printed earlier may move. Source 63's limit supersedes source 60's bracket, so source 63 comes first; source 60's section is a second route with its own connected-limit theorem. Source 63's Section 2 also fixes the conventions of its Section 6 (Section IV.67).

What is proved. Source 63: for real centered f on any abelian group of odd order $N \geq 3$, $\|f\|_{U^2}^4 \leq (1/3 - 4/(3N^2) + 1/N^3) \|f\|_{\infty}^4$, with equality on $\mathbb{Z}/N\mathbb{Z}$, so this is the best bound depending only on N , and $1/3$ is sharp uniformly (Theorems III.96.2, III.96.4; a compact-group version, Corollary III.96.6); then $\lim_d c_d^{\text{odd}} = 1/3$ with $c_d^{\text{odd}} \leq 1/3 + 10 \cdot 2^{-\lfloor d/8 \rfloor}$ for $d \geq 80$ (Theorem III.97.1), an explicit finite-parameter bound and threshold (Theorem III.97.5, Corollary III.97.6), the limit $b(G)$ on each fixed group (Corollary III.97.7), and for the quartic cube coefficient $A_d = P_d c_d^{\text{odd}} \sim 6^d/24$ (Theorem III.98.5). *Source 60:* $\lim_d c_d^{\text{odd}} = \beta_{\text{odd}}$, the bounded Fourier extremum, with $1/3 \leq \beta_{\text{odd}} \leq 1/2 - 8/\pi^2 + 64/\pi^4 = 0.3464533951$ (Theorem III.99.1), and the joint limit $1/3$ when the least prime divisor also grows, through the Riesz–Sobolev inequality of Christ and Iliopoulou (Theorem III.99.2). Source 63's inequality gives $\beta_{\text{odd}} = 1/3$, so source 60's limit theorem now reads $\lim_d c_d^{\text{odd}} = 1/3$ and its bracket is superseded; source 60's statements are correct as printed (note after Theorem III.96.4). This answers the limit questions of sources 28 and 29 and source 60's Research questions 1–2 and the first half of 3 (item III.1). At fixed d the bracket of the report stands: $0.5587725 \leq c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ (the batch-121 certificate and source 49), and source 63's rate is compatible with it; neither source determines c_d^{odd} at any fixed $d \geq 4$, as both say. Sources placed during this write record $c_4^{\text{odd}} \leq 0.5868716$ (source 96) and a rate $O(d2^{-d/2})$ (source 93); they are not checked here (note after Theorem III.97.1). Source 63 does not cite source 60, nor conversely; both pins precede the other's arrival.

Relation to the formal statements (checked at HEAD fb9602e55). `lemma_3_10` (Sections01_03.lean, line 340) has the exact companion `lemma_3_10_holds` (Proofs03CubeUpper.lean, line 82); no declaration concerns c_d^{odd} . Nothing in this chapter is formalized.

III.95 Conventions and standard tools

Source 63, delivered Section 2.

Unless specified otherwise, G is a finite abelian group written additively, with normalized averaging $\mathbb{E}_{x \in G} = |G|^{-1} \sum_{x \in G}$. A function is *centered* if its mean is zero. We use $\mathcal{C}z = \bar{z}$ and $\omega \cdot h = \sum_i \omega_i h_i$. All logarithms without a subscript are natural. A factorial with argument zero is 1.

For $r \geq 1$, the normalized Gowers expression is

$$\|f\|_{U^r(G)}^{2^r} = \mathbb{E}_{x, h_1, \dots, h_r} \prod_{\omega \in \{0,1\}^r} \mathcal{C}^{|\omega|} f(x + \omega \cdot h). \quad (\text{III.95.1})$$

For $r = 1$, this is $|\mathbb{E}f|^2$, and for $r \geq 2$ its 2^r th root is a norm. With $\Delta_h f(x) = f(x) \overline{f(x+h)}$, the

recursion is

$$\|f\|_{U^{r+1}}^{2^{r+1}} = \mathbb{E}_h \|\Delta_h f\|_{U^r}^{2^r}. \quad (\text{III.95.2})$$

For real f , all conjugations may be omitted, but the resulting individual cube products can still be negative.

We use the mixed Gowers–Cauchy–Schwarz inequality

$$\left| \mathbb{E}_{x,h} \prod_{\omega \in \{0,1\}^r} \mathcal{C}^{|\omega|} f_\omega(x + \omega \cdot h) \right| \leq \prod_{\omega \in \{0,1\}^r} \|f_\omega\|_{U^r}. \quad (\text{III.95.3})$$

This is the normalized form of Gowers’ Lemma 3.8. Iterated Cauchy–Schwarz in the cube coordinates proves it; expanding the cube of $f + g$ and applying (III.95.3) gives the norm triangle inequality of his Lemma 3.9. The same identities yield

$$|\mathbb{E} f| \leq \|f\|_{U^2} \leq \|f\|_{U^3} \leq \cdots \leq \|f\|_\infty. \quad (\text{III.95.4})$$

These facts hold on every finite abelian group. Gowers’ original presentation and Tao’s account provide the standard proofs [1, 257]. We use them as established tools; the new estimates appear in the subsequent sections.

For a character $\chi \in \widehat{G}$, let $\widehat{f}(\chi) = \mathbb{E}_x f(x) \overline{\chi(x)}$. Fourier inversion, orthogonality, and normalized convolution give

$$f(x) = \sum_{\chi \in \widehat{G}} \widehat{f}(\chi) \chi(x), \quad \mathbb{E} |f|^2 = \sum_{\chi} |\widehat{f}(\chi)|^2, \quad (\text{III.95.5})$$

$$(f * \mu)(x) = \mathbb{E}_y f(x - y) \mu(y), \quad \widehat{f * \mu}(\chi) = \widehat{f}(\chi) \widehat{\mu}(\chi). \quad (\text{III.95.6})$$

In particular,

$$\|f\|_{U^2}^4 = \sum_{\chi} |\widehat{f}(\chi)|^4. \quad (\text{III.95.7})$$

For real f , $\widehat{f}(\chi^{-1}) = \overline{\widehat{f}(\chi)}$. If G has odd order, each nontrivial character is distinct from its inverse. These conjugate pairs are central to the uniform Fourier-tail argument.

The local phase-removal section deliberately uses unnormalized cube sums $C_r(f) = |G|^{r+1} \|f\|_{U^r}^{2^r}$, because its input Fourier sums and progression sizes are unnormalized. The conversion is always displayed. The letter d denotes a norm order outside that section; inside the local problem it also denotes the common difference of the input progressions, as in Gowers’ statement.

III.96 A sharp bound for bounded real functions

Source 63, delivered Section 3.

The first ingredient is a pointwise constraint on autocorrelations. It uses both real-valuedness and the bounded range of the function. The constraint is stronger than positive definiteness alone.

Lemma III.96.1 (The doubling constraint). *Let G be a finite abelian group and let $f : G \rightarrow [-1, 1]$ be real. For*

$$R_f(h) = \mathbb{E}_{x \in G} f(x) f(x + h)$$

one has

$$1 + R_f(2h) \geq 2|R_f(h)| \quad (h \in G). \quad (\text{III.96.1})$$

Proof. For real $u, v, w \in [-1, 1]$,

$$uv + vw - uw \leq 1, \quad -uv - vw - uw \leq 1. \quad (\text{III.96.2})$$

The first expression is affine separately in each variable, so its maximum on the cube is attained at a vertex. At a vertex with $u = w$, it is $2uv - 1 \leq 1$; at a vertex with $u = -w$, it equals 1. The second inequality follows by replacing v by $-v$. Apply these inequalities to $(u, v, w) = (f(x), f(x+h), f(x+2h))$ and average over x . Translation invariance gives $\pm 2R_f(h) - R_f(2h) \leq 1$, which is the assertion. $\square$

Theorem III.96.2 (The sharp bounded centered inequality). *Let G be a finite abelian group of odd order and let $f : G \rightarrow \mathbb{R}$ have mean zero. Then*

$$\|f\|_{U^2(G)}^4 \leq \frac{1}{3} \|f\|_\infty^4. \quad (\text{III.96.3})$$

The constant $1/3$ cannot be decreased uniformly over such groups.

Proof. The zero function is immediate; otherwise scale to $\|f\|_\infty = 1$. Write $R = R_f$ and $S = \mathbb{E}_h R(h)^2 = \|f\|_{U^2}^4$. Both sides of (III.96.1) are nonnegative, so

$$(1 + R(2h))^2 \geq 4R(h)^2.$$

Doubling is a bijection on a finite group of odd order. Moreover, $\mathbb{E}_h R(h) = (\mathbb{E} f)^2 = 0$. Averaging the last display therefore gives $1 + S \geq 4S$, and hence $S \leq 1/3$. Sharpness follows from the exact cyclic example in Theorem III.96.4 below, by letting the odd order tend to infinity. $\square$

Remark III.96.3 (The precise role of odd order). On an even-order abelian group there is a nontrivial character with values in $\{-1, 1\}$. It is real and centered, and both its supremum norm and its U^d norm are 1 for every $d \geq 2$. Thus neither (III.96.3) nor the limiting odd-order comparison proved below extends unchanged to all finite groups. The proof of Theorem III.96.2 uses odd order exactly to make averaging over h and $2h$ agree.

III.96.1 An exact finite-order correction

The value $1/3$ is a supremum over group orders. At each fixed odd order there is a sharp strict correction. The identity behind the proof makes it possible to determine that correction without classifying all autocorrelation sequences.

Theorem III.96.4 (Sharp universal bound at an odd order). *Let $N \geq 3$ be odd, let G be an abelian group of order N , and let $f : G \rightarrow \mathbb{R}$ have mean zero. Then*

$$\|f\|_{U^2(G)}^4 \leq \left(\frac{1}{3} - \frac{4}{3N^2} + \frac{1}{N^3} \right) \|f\|_\infty^4. \quad (\text{III.96.4})$$

For every odd N , equality is attained on $G = \mathbb{Z}/N\mathbb{Z}$. Consequently this is the best bound depending only on the order N .

Proof. Normalize as before, put $r = \mathbb{E} f^2$, and define the nonnegative slack

$$D_f(h) = (1 + R_f(2h))^2 - 4R_f(h)^2.$$

The averaging argument gives the exact identity

$$1 - 3\|f\|_{U^2}^4 = \mathbb{E}_h D_f(h). \quad (\text{III.96.5})$$

At the identity element,

$$D_f(0) = (1 - r)(1 + 3r). \quad (\text{III.96.6})$$

We first establish $r \leq 1 - 1/N$. Consider the compact polytope

$$\mathcal{P}_N = \{(t_1, \dots, t_N) \in [-1, 1]^N : t_1 + \dots + t_N = 0\}.$$

At an extreme point, at most one coordinate lies strictly between -1 and 1 : two such coordinates could be perturbed in opposite directions. The other $N - 1$ coordinates have values ± 1 and an even integer sum. Its negative must lie in $[-1, 1]$, so the remaining coordinate is 0 and the sum of the others is 0 . The convex function $\sum t_i^2$ has a maximum at an extreme point, where its value is $N - 1$. This proves the claim.

The concave function $(1 - r)(1 + 3r)$ takes its minimum on $[0, 1 - 1/N]$ at an endpoint. Its endpoint values are 1 and $4/N - 3/N^2$, and the latter is at most 1 for $N \geq 3$. Thus, by (III.96.5)–(III.96.6),

$$1 - 3\|f\|_{U^2}^4 \geq \frac{D_f(0)}{N} \geq \frac{4}{N^2} - \frac{3}{N^3}.$$

This is (III.96.4).

For attainment, write $N = 2m + 1$ and define on $\mathbb{Z}/N\mathbb{Z}$

$$f_N(0) = 0, \quad f_N(j) = 1 \ (1 \leq j \leq m), \quad f_N(j) = -1 \ (m + 1 \leq j \leq 2m). \quad (\text{III.96.7})$$

This function is centered and has supremum norm 1 . For $1 \leq h \leq m$, two terms in $\sum_x f_N(x)f_N(x+h)$ vanish, and exactly $2h - 1$ of the remaining terms are negative. Consequently

$$R_{f_N}(0) = 1 - \frac{1}{N}, \quad R_{f_N}(h) = 1 - \frac{4h}{N} \ (1 \leq h \leq m), \quad R_{f_N}(-h) = R_{f_N}(h).$$

The formulas for the sums of the first m integers and their squares now give

$$\|f_N\|_{U^2}^4 = \frac{1}{N} \left[\left(1 - \frac{1}{N}\right)^2 + 2 \sum_{h=1}^m \left(1 - \frac{4h}{N}\right)^2 \right] = \frac{1}{3} - \frac{4}{3N^2} + \frac{1}{N^3}.$$

This proves sharpness at every odd order. □

[write] 7 October 2026; sources 60–69. *Checked at the write* (exact rationals). $1/3 - 4/(3N^2) + 1/N^3 = (N^3 - 4N + 3)/(3N^3)$, which is source 60's conjectured cyclic value, and it equals source 60's exhaustive maxima $2/9, 36/125, 106/343, 232/729, 430/1331, 716/2197, 1106/3375, 1616/4913$ for $N = 3, 5, \dots, 17$; on $\mathbb{Z}_3 \times \mathbb{Z}_3$ the maximum $56/243$ is below the bound $232/729$, so equality depends on the group, not only on N . *Consequences in this report:* source 60's β_{odd} equals $1/3$, so its bracket $[1/3, 0.3464533951]$ is superseded, and its Research questions V.226.1 and V.226.3 (first part) are answered (notes there).

Remark III.96.5 (What the finite theorem classifies). Theorem III.96.4 determines the optimal bound uniform over groups of the specified order and supplies exact extremizers. It does not assert that every abelian group of that order attains the bound, or classify all equality cases on a fixed group. Those are distinct structural questions. The finite correction is also separate from the U^d denominator problem: controlling $\|f\|_\infty$ by $\|f\|_{U^d}$ with an error uniform over arbitrary groups requires the arguments of Section III.97.

III.96.2 Compact groups and a defect identity

Corollary III.96.6 (A compact-group version). *Let G be a compact abelian group such that $x \mapsto 2x$ is surjective. With normalized Haar measure and the essential supremum norm, every bounded measurable real centered function satisfies*

$$\|f\|_{U^2(G)}^4 \leq \frac{1}{3} \|f\|_\infty^4.$$

On $G = \mathbb{R}/\mathbb{Z}$ equality is attained by $f(x) = \text{sgn}(\sin(2\pi x))$, with arbitrary values at its two zeros.

Proof. The pushforward of normalized Haar measure under a surjective continuous homomorphism is normalized Haar measure. Therefore the proof of Theorem III.96.2, including Fubini's theorem for the correlation average, applies verbatim. For the displayed square wave the autocorrelation is $R(h) = 1 - 4\|h\|_{\mathbb{R}/\mathbb{Z}}$. Thus

$$\|f\|_{U^2}^4 = 2 \int_0^{1/2} (1 - 4t)^2 dt = \frac{1}{3}.$$

□

The proof gives more than an upper bound. For a centered real 1-bounded function on either class of groups,

$$1 - 3\|f\|_{U^2}^4 = \mathbb{E}_h(1 + R_f(2h) - 2|R_f(h)|)(1 + R_f(2h) + 2|R_f(h)|). \quad (\text{III.96.8})$$

Both factors are nonnegative. Hence a function close to the sharp bound must nearly saturate the doubling constraint for most increments on which the second factor is bounded away from zero. For example, if $1 - 3\|f\|_{U^2}^4 \leq \varepsilon$, then for all $a, b > 0$,

$$\Pr(1 + R_f(2h) - 2|R_f(h)| \geq a, \quad 1 + R_f(2h) + 2|R_f(h)| \geq b) \leq \frac{\varepsilon}{ab}.$$

This is an exact analytic starting point for stability. It does not by itself imply proximity of f to a cyclic square wave.

III.97 The exact high-order odd-group comparison constant

Source 63, delivered Section 4.

For each integer $d \geq 2$, define

$$c_d^{\text{odd}} := \sup_{\substack{G \text{ finite abelian, } |G| \text{ odd} \\ f: G \rightarrow \mathbb{R}, \mathbb{E}f=0, f \neq 0}} \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4}. \quad (\text{III.97.1})$$

The supremum includes the group. Consequently, convergence of Gowers norms for a fixed finite group would not by itself determine the limit of these constants. The purpose of this section is to establish the required uniformity.

Theorem III.97.1 (Exact limiting constant and a uniform rate). *The constants in (III.97.1) satisfy*

$$\lim_{d \rightarrow \infty} c_d^{\text{odd}} = \frac{1}{3}. \quad (\text{III.97.2})$$

More explicitly, for every integer $d \geq 80$,

$$\frac{1}{3} \leq c_d^{\text{odd}} \leq \frac{1}{3} + 10 \cdot 2^{-\lfloor d/8 \rfloor}. \quad (\text{III.97.3})$$

Neither the exponent $1/8$ nor the coefficient 10 in this rate is asserted optimal.

[write] 7 October 2026; sources 60–69. *This answers* the limit questions of sources 28 and 29 (Question after Proposition III.69.2 and source 28's problem on the asymptotic constants; notes there) and source 60's Research question V.226.2 for c_d^{odd} . At fixed d it is compatible with the report's bracket $0.5587725 \leq c_4^{\text{odd}} \leq 0.6094183$ (sources 28, 49 and the batch-121 certificate) and with source 49's $c_5^{\text{odd}} \leq 29/50$; no fixed- d value is claimed. *Forward pointers* (from placement records, not checked at this write): source 93 records a rate $c_d(F) - D(F) = O(d2^{-d/2})$ uniformly over odd-order families, which would supersede the rate here; source 96 records $\|f\|_{U^4}^{16} \geq (843/100) \|f\|_{U^2}^{16}$, that is $c_4^{\text{odd}} \leq 0.5868716$, below source 49's 0.6094183. Both are for the writes that add them.

The earlier repository report established $1/3 \leq \lim_d c_d^{\text{odd}} \leq 1/2$ by an interval construction and central-binomial autocorrelation moments [258]. Theorem III.97.1 identifies the missing endpoint. Its proof combines the bounded-function inequality of Theorem III.96.2 with three uniform estimates. High even moments control the tail of the Fourier fourth powers. A positive convolution kernel retains the important coefficients and has finite Fourier support. Finally, small increments in a Bohr set recover the supremum norm of this smoothed function from a sufficiently high Gowers norm. Every estimate below is independent of the order and exponent of the finite group.

III.97.1 Uniform control of the Fourier tail

We use normalized Fourier coefficients $\widehat{f}(\chi) = \mathbb{E}_x f(x) \overline{\chi(x)}$. For real f , its autocorrelation

$$R_f(h) = \mathbb{E}_x f(x) f(x+h) = \sum_{\chi \in \widehat{G}} |\widehat{f}(\chi)|^2 \chi(h) \quad (\text{III.97.4})$$

has nonnegative Fourier coefficients. This positivity permits us to retain selected words in each even-moment expansion. The following estimate specifically controls mass spread over many small coefficients; a bound on the largest coefficient would not suffice.

Lemma III.97.2 (Fourier-tail estimate). *Suppose G has odd order, $f : G \rightarrow \mathbb{R}$ is centered, and $\|f\|_{U^d} = 1$. Choose one representative ξ_j from each nonzero conjugate Fourier pair, ordered so that $|\widehat{f}(\xi_j)|$ decreases, and put*

$$v_j = 2|\widehat{f}(\xi_j)|^4, \quad S = \sum_j v_j = \|f\|_{U^2}^4 \leq 1.$$

For integers $1 \leq r \leq 2^{d-2}$ and $K \geq 1$, let

$$G_r = (2r-1)!! = \frac{(2r)!}{2^r r!}, \quad \tau(r, K) = \max \left\{ \frac{r(r-1)}{K+1}, \left(\frac{2}{G_r} \right)^{1/r} \right\}. \quad (\text{III.97.5})$$

Then

$$\sum_{j>K} v_j \leq \tau(r, K). \quad (\text{III.97.6})$$

An empty tail is interpreted as zero. In particular, for $r \geq 16$,

$$\tau(r, r^3) \leq \frac{3}{r}. \quad (\text{III.97.7})$$

Proof. Put $n = 2^{d-2}$. The derivative recursion and the inequality $\|g\|_{U^{d-1}} \geq |\mathbb{E}g|$ give

$$1 = \|f\|_{U^d}^{2^d} \geq \mathbb{E}_h |R_f(h)|^{2n}.$$

Monotonicity of probability-space L^p norms therefore gives $\mathbb{E}R_f^{2r} \leq 1$ for $r \leq n$.

Write $\lambda_j = |\widehat{f}(\xi_j)|^2$, so $R_f = \sum_j \lambda_j (\xi_j + \xi_j^{-1})$ and $v_j = 2\lambda_j^2$. Each surviving word in the expansion of $\mathbb{E}R_f^{2r}$ has a nonnegative coefficient. Retain only words using r distinct tail pairs, each once in each orientation. For a fixed set of pairs their coefficient is $(2r)!$. Thus, with e_r denoting the elementary symmetric polynomial,

$$1 \geq \mathbb{E}R_f^{2r} \geq \frac{(2r)!}{2^r} e_r(v_{K+1}, v_{K+2}, \dots). \quad (\text{III.97.8})$$

Additional character relations can only add nonnegative terms, so no independence or restriction on odd torsion has entered this argument.

Set $v = \sum_{j>K} v_j$. If $v = 0$ there is nothing to prove. Draw r independent tail indices with probabilities v_j/v . A union bound for collisions gives

$$\frac{r! e_r(v_{K+1}, \dots)}{v^r} \geq 1 - \binom{r}{2} \sum_{j>K} \left(\frac{v_j}{v}\right)^2 \geq 1 - \binom{r}{2} \frac{v_{K+1}}{v}.$$

Since the v_j decrease and their sum is at most one, $v_{K+1} \leq 1/(K+1)$. If $v \geq r(r-1)/(K+1)$, the last display is at least $1/2$. Substitution in (III.97.8) yields $1 \geq G_r v^r/2$. If that hypothesis fails, the other term of (III.97.5) already bounds v . This proves (III.97.6).

For $K = r^3$, the first term of (III.97.5) is at most $1/r$. The elementary integral estimate

$$\log(r!) \geq \int_1^r \log x \, dx \geq r \log r - r$$

and $G_r \geq r!$ give $(2/G_r)^{1/r} \leq 2^{1/r} e/r < 3/r$ for $r \geq 16$. For example, $e < 11/4$ and $2^{1/16} < 12/11$ justify the last strict inequality without any asymptotic estimate. $\square$

III.97.2 A positive smoothing kernel with a finite spectrum

We next preserve a prescribed collection of Fourier coefficients. Using a product of trigonometric sums, rather than assuming that the chosen characters are independent, is essential on arbitrary finite groups.

Lemma III.97.3 (Positive finite-spectrum smoothing). *Let $\xi_1, \dots, \xi_\ell$ be characters of a finite abelian group, where $1 \leq \ell \leq K$, and let $J \geq 1$ be an integer. Define*

$$P(x) = \prod_{j=1}^{\ell} \left(\sum_{a=0}^J \xi_j(x)^a \right), \quad \mu(x) = \frac{|P(x)|^2}{\mathbb{E}|P|^2}. \quad (\text{III.97.9})$$

Then $\mu \geq 0$, $\mathbb{E}\mu = 1$, and its Fourier support is contained in

$$\Gamma = \left\{ \prod_{j=1}^{\ell} \xi_j^{a_j} : -J \leq a_j \leq J \right\}, \quad |\Gamma| \leq (2J+1)^K. \quad (\text{III.97.10})$$

For each selected character,

$$\widehat{\mu}(\xi_j) \geq \frac{J}{J+1}. \quad (\text{III.97.11})$$

All Fourier coefficients of μ are nonnegative real numbers. For $g = f * \mu$, normalized convolution satisfies

$$\|g\|_{U^d} \leq \|f\|_{U^d} \quad (d \geq 2). \quad (\text{III.97.12})$$

If f is real and centered, so is g .

Proof. The Fourier coefficients of P are nonnegative integers, including when different exponent vectors represent the same character. Thus $\mathbb{E}|P|^2 > 0$, and every coefficient of $|P|^2$ is nonnegative real. Its spectrum is contained in (III.97.10).

Fix j , and let P_j be the product in (III.97.9) with its j th factor removed. Every correlation $\langle P_j, \chi P_j \rangle$ is nonnegative real, by the same Fourier-coefficient observation. Consequently,

$$\begin{aligned} \|P\|_2^2 &\geq (J+1) \|P_j\|_2^2, \\ \|P - \xi_j P\|_2^2 &= \|(1 - \xi_j^{J+1})P_j\|_2^2 \leq 2 \|P_j\|_2^2. \end{aligned}$$

Here $\|P\|_2$ denotes the normalized L^2 norm. Since

$$1 - \hat{\mu}(\xi_j) = \frac{\|P - \xi_j P\|_2^2}{2 \|P\|_2^2},$$

the multiplier bound follows. This computation explicitly retains all aliases caused by relations among the selected characters. Finally, convolution with μ is an average of translates with nonnegative weights. The triangle inequality and translation invariance of U^d prove (III.97.12). The assertions about reality and the mean follow directly from convolution. $\square$

III.97.3 Uniform recovery of a supremum norm

The next lemma supplies the uniformity needed when the order of the group varies. The point of using a square of a parallel-face average is that the full cube integrand may have either sign.

Lemma III.97.4 (Finite-spectrum recovery). *Suppose $g : G \rightarrow \mathbb{R}$ has Fourier support contained in (III.97.10) and satisfies $\|g\|_{U^d} \leq 1$, where $d \geq 2$. Put*

$$M = (2J+1)^K, \quad L = 2\pi M K J. \quad (\text{III.97.13})$$

For a real parameter $0 < \eta \leq 1/16$, define

$$\beta = \left(\frac{\eta}{2L}\right)^K, \quad \gamma = \left(\frac{\eta}{2(d-1)L}\right)^K, \quad s = 2^{d-1}, \quad (\text{III.97.14})$$

$$F = \gamma^{(d-1)/2^d} (\beta/2)^{2/2^d}. \quad (\text{III.97.15})$$

If

$$3^{-s} \leq \beta/2, \quad (\text{III.97.16})$$

then

$$\|g\|_\infty \leq 2\eta + F^{-1}. \quad (\text{III.97.17})$$

All quantities on the right depend only on K, J, η, d .

Proof. For $z \in \mathbb{R}/\mathbb{Z}$, write $\|z\|_{\mathbb{R}/\mathbb{Z}}$ for its distance to zero, and put

$$\rho(y) = \max_{1 \leq j \leq \ell} \|\arg \xi_j(y)/(2\pi)\|_{\mathbb{R}/\mathbb{Z}}, \quad B(t) = \{y : \rho(y) \leq t\}.$$

The argument of each character is interpreted as an element of $\mathbb{R}/\mathbb{Z}$, so no choice of a real branch is involved. For $0 < t \leq 1$,

$$\frac{|B(t)|}{|G|} \geq (t/2)^K. \quad (\text{III.97.18})$$

Indeed, partition each of the ℓ circle coordinates into $\lceil 1/t \rceil$ equal half-open intervals. Some product cell contains at least $\lceil 1/t \rceil^{-\ell}$ of the image measure. Translate its preimage by one of its points; the resulting differences lie in $B(t)$. Now use $\lceil 1/t \rceil \leq 2/t$ and $\ell \leq K$.

Every Fourier coefficient of g has modulus at most $\|g\|_{U^2} \leq 1$. Select for each member of Γ one representation in (III.97.10). The triangle inequality on the circle gives

$$|g(x+y) - g(x)| \leq L\rho(y). \quad (\text{III.97.19})$$

There are at most M coefficients, and each chosen frequency has total word length at most KJ , which proves this estimate with the value of L in (III.97.13).

Let $A = \|g\|_{\infty}$. If $A < 1/2$, (III.97.17) is immediate because $F \leq 1$. Otherwise, choose x_0 with $|g(x_0)| = A$. On $x_0 + B(\eta/L)$ we have $|g(x)| \geq A - \eta$; this set has measure at least β by (III.97.18).

Restrict each of $h_1, \dots, h_{d-1}$ to $B(\eta/((d-1)L))$. The measure of each allowed increment set is at least γ . For every cube vertex ω , the subadditivity of ρ and (III.97.19) give

$$|g(x + \omega \cdot h) - g(x)| \leq \eta.$$

If $|g(x)| \geq \eta$, the s factors share the same weak sign. Since s is even, their product is nonnegative and is at least $(|g(x)| - \eta)^s$. If $|g(x)| < \eta$, every factor has modulus at most 2η , so the product is at least $-(2\eta)^s$. Thus the entire parallel-face average, including the region away from x_0 , satisfies

$$\begin{aligned} A_h &:= \mathbb{E}_x \prod_{\omega \in \{0,1\}^{d-1}} g(x + \omega \cdot h) \\ &\geq \beta(A - 2\eta)^s - (2\eta)^s \geq \frac{\beta}{2}(A - 2\eta)^s. \end{aligned}$$

For the last inequality, $A \geq 1/2$ and $\eta \leq 1/16$ imply $2\eta/(A - 2\eta) \leq 1/3$, and (III.97.16) applies.

The exact parallel-face identity for a real function is

$$\|g\|_{U^d}^{2^d} = \mathbb{E}_{h_1, \dots, h_{d-1}} |A_h|^2.$$

Its summands are nonnegative. Retaining the restricted increments therefore gives

$$1 \geq \|g\|_{U^d}^{2^d} \geq \gamma^{d-1}(\beta/2)^2(A - 2\eta)^{2^d}.$$

Taking 2^d th roots proves (III.97.17). □

III.97.4 A finite-parameter comparison theorem

Theorem III.97.5 (Explicit finite-parameter bound). *Let $d \geq 2$, $1 \leq r \leq 2^{d-2}$, $K, J \geq 1$ be integers, and $0 < \eta \leq 1/16$. Define τ by (III.97.5), and M, L, β, γ, F by (III.97.13)–(III.97.15). If (III.97.16) holds, then*

$$c_d^{\text{odd}} \leq \frac{1}{3} \left(1 + \frac{1}{J}\right)^4 (2\eta + F^{-1})^4 + \tau(r, K). \quad (\text{III.97.20})$$

This is a finite inequality, with no limiting qualification.

Proof. Take any admissible G, f , and use homogeneity to normalize $\|f\|_{U^d} = 1$. Set $S = \|f\|_{U^2}^4$ and select the first K nonzero Fourier pairs as in Lemma III.97.2. If fewer than K pairs occur, select all of them and use their actual number $\ell \leq K$ in the kernel. All complexity bounds remain valid with K in place of ℓ .

Let $g = f * \mu$, using Lemma III.97.3. Its multiplier bound and the Fourier formula for U^2 imply

$$\|g\|_{U^2}^4 \geq \left(\frac{J}{J+1}\right)^4 \sum_{j \leq K} 2|\widehat{f}(\xi_j)|^4 \geq \left(\frac{J}{J+1}\right)^4 (S - \tau(r, K)).$$

The function g is real centered and satisfies $\|g\|_{U^d} \leq 1$. Apply Theorem III.96.2 and then Lemma III.97.4 to obtain

$$\|g\|_{U^2}^4 \leq \frac{1}{3} \|g\|_{\infty}^4 \leq \frac{1}{3} (2\eta + F^{-1})^4.$$

Rearrange and take the supremum over G, f . □

Proof of the limiting assertion in Theorem III.97.1. First fix a sufficiently large integer r , and take $K = r^3$. Then $\tau(r, K) \leq 3/r$. Next choose J large and $\eta > 0$ small. These parameters are fixed independently of the group. As $d \rightarrow \infty$, (III.97.16) eventually holds, and

$$-\log F = O_{K,J,\eta} \left(\frac{d \log d}{2^d} \right) \rightarrow 0.$$

Theorem III.97.5 therefore implies

$$\limsup_{d \rightarrow \infty} c_d^{\text{odd}} \leq \frac{1}{3} (1 + J^{-1})^4 (1 + 2\eta)^4 + \frac{3}{r}.$$

Let $J, r \rightarrow \infty$ and $\eta \downarrow 0$. This proves the upper bound $1/3$, without interchanging a fixed-group limit with a supremum.

For completeness, recall the lower construction from the repository [258]. On $G = \mathbb{Z}/(2n+1)\mathbb{Z}$, let $A = \{0, \dots, n-1\}$, $\delta = n/(2n+1)$, and $f = 2(\mathbf{1}_A - \delta)$. There is no wraparound in sums of two members of A , so its unnormalized additive energy is

$$E(A) = n^2 + 2 \sum_{j=1}^{n-1} j^2 = \frac{2n^3 + n}{3}.$$

Fourier orthogonality and direct evaluation give

$$\|f\|_{U^2}^4 = 16 \left(\frac{E(A)}{(2n+1)^3} - \delta^4 \right), \quad \|f\|_{\infty} = \frac{2(n+1)}{2n+1}.$$

Since $\|f\|_{U^d} \leq \|f\|_\infty$, for every $d \geq 2$,

$$\frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4} \geq \frac{n(n^2 + n + 1)}{3(n+1)^3} \longrightarrow \frac{1}{3}.$$

Hence $c_d^{\text{odd}} \geq 1/3$ for every d , completing the proof of (III.97.2). $\square$

III.97.5 An explicit convergence rate

The finite theorem allows many parameter choices. The following deliberately simple choice proves (III.97.3); optimizing this construction is a separate question.

Proof of the quantitative assertion in Theorem III.97.1. Let $k \geq 10$ be an integer and put

$$d = 8k, \quad r = 2^k, \quad K = r^3, \quad J = r, \quad \eta = r^{-1}. \quad (\text{III.97.21})$$

Then $r \leq 2^{d-2}$, $\eta \leq 1/16$, and $\tau(r, K) \leq 3/r$. Write

$$H = \log(2L/\eta), \quad t = -\log F.$$

Since $\log(2\pi) < 3$, $\log(2r+1) \leq k+2$, and $\log r \leq k$, the explicit definitions yield

$$H \leq r^3(k+2) + 5k + 4, \quad (\text{III.97.22})$$

$$t = \frac{(8k+1)r^3H + (8k-1)r^3\log(8k-1) + 2\log 2}{r^8}. \quad (\text{III.97.23})$$

For every $k \geq 10$,

$$\begin{aligned} (8k+1)(5k+4) + (8k-1)\log(8k-1) + 2 &\leq 2^{3k}, \\ (8k+1)(k+2) + 1 &\leq 2^k. \end{aligned}$$

These inequalities can be checked elementarily. In the first, replace $\log(8k-1)$ by $8k-1$ and compare the resulting quadratic polynomial with 2^{3k} , starting at $k=10$. For the second, the initial comparison is $973 \leq 1024$, and the quadratic polynomial grows by less than a factor of two at each subsequent step. Substitution in (III.97.22)–(III.97.23) gives

$$0 \leq t \leq \frac{(8k+1)(k+2) + 1}{r^2} \leq \frac{1}{r}. \quad (\text{III.97.24})$$

We also check the activation condition. From (III.97.22), $H \leq r^3(k+3)$, and hence

$$-\log(\beta/2) = r^3H + \log 2 \leq r^6(k+4).$$

On the other hand, $2^{d-1}\log 3 > r^8/2 \geq r^6(k+4)$. Thus (III.97.16) holds.

For $0 \leq x \leq 1$, $e^x \leq 1 + x + x^2$. Since $r \geq 1024$,

$$\begin{aligned} \left(1 + \frac{1}{r}\right) \left(\frac{2}{r} + e^{1/r}\right) &\leq 1 + \frac{4}{r} + \frac{4}{r^2} + \frac{1}{r^3} \leq 1 + \frac{5}{r}, \\ \left(1 + \frac{5}{r}\right)^4 &\leq 1 + \frac{21}{r}. \end{aligned}$$

Theorem III.97.5 and (III.97.24) now imply

$$c_{8k}^{\text{odd}} \leq \frac{1}{3} \left(1 + \frac{21}{r} \right) + \frac{3}{r} = \frac{1}{3} + \frac{10}{r}.$$

Finally, norm monotonicity makes c_d^{odd} nonincreasing in d . For arbitrary $d \geq 80$, take $k = \lfloor d/8 \rfloor$ and use $c_d^{\text{odd}} \leq c_{8k}^{\text{odd}}$. $\square$

Corollary III.97.6 (A constructive dimension threshold). *For $\varepsilon > 0$, set*

$$d_\varepsilon = 8 \max \left\{ 10, \left\lceil \log_2 \frac{10}{\varepsilon} \right\rceil \right\}.$$

For every $d \geq d_\varepsilon$, every finite odd-order abelian group, and every real centered f on that group,

$$\|f\|_{U^2}^4 \leq \left(\frac{1}{3} + \varepsilon \right) \|f\|_{U^d}^4.$$

The threshold is uniform in both the function and the group. It therefore supplies the quantifier order required by the original extremal question. It does not determine the optimal constant at any fixed $d \geq 4$, and the parameter choices in (III.97.21) should not be read as evidence for an optimal exponential rate.

Corollary III.97.7 (Exact limits on a fixed finite group). *For a fixed finite abelian group G of odd order $N \geq 3$, put*

$$b(G) = \sup_{\substack{f: G \rightarrow \mathbb{R}, \mathbb{E}f=0 \\ f \neq 0}} \frac{\|f\|_{U^2}^4}{\|f\|_\infty^4}, \quad c_d(G) = \sup_{\substack{f: G \rightarrow \mathbb{R}, \mathbb{E}f=0 \\ f \neq 0}} \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4}.$$

Then, for every $d \geq 2$,

$$b(G) \leq c_d(G) \leq b(G) N^{4(d+1)/2^d}, \quad \lim_{d \rightarrow \infty} c_d(G) = b(G). \quad (\text{III.97.25})$$

In particular, on every odd cyclic group,

$$\lim_{d \rightarrow \infty} c_d(\mathbb{Z}/N\mathbb{Z}) = \frac{1}{3} - \frac{4}{3N^2} + \frac{1}{N^3}. \quad (\text{III.97.26})$$

Proof. The inequality $\|f\|_{U^d} \leq \|f\|_\infty$ gives the lower bound in (III.97.25). For the upper bound, use the parallel-face square identity and retain only the increments $h_1 = \dots = h_{d-1} = 0$. Since 2^{d-1} is even,

$$\|f\|_{U^d}^{2^d} \geq N^{-(d-1)} \left(\mathbb{E}_x f(x)^{2^{d-1}} \right)^2 \geq N^{-(d+1)} \|f\|_\infty^{2^d}.$$


Thus $\|f\|_\infty^4 \leq N^{4(d+1)/2^d} \|f\|_{U^d}^4$, proving the upper bound after taking the supremum. The multiplier tends to one for fixed N . The cyclic value of $b(G)$ is exactly Theorem III.96.4. $\square$

The dependence on N in (III.97.25) explains why this short fixed-group argument does not prove Theorem III.97.1: its error can be arbitrarily large when the group order varies with d . The Fourier-tail and smoothing arguments remove that dependence.

The accompanying `code/verify_norm_limit.py` checks the distinct-pair factorials and collision estimates by exact integer arithmetic on six finite groups, checks the smoothing multipliers in examples with character aliases, and evaluates the finite parameter expressions with 100-digit arithmetic. These finite checks supplement the proofs; the decimal evaluation is not a certified interval-arithmetic calculation.

d	Bound from (III.97.20)	Simple bound from (III.97.3)
64	0.36328010	—
80	0.33938108	0.34309896
96	0.33471262	0.33577474
112	0.33366635	0.33394369
128	0.33341558	0.33348592
160	0.33333845	0.33334287

Table III.97.1: Approximate values of the proved upper-bound expressions. The middle column uses (III.97.21), with (III.97.20) evaluated directly; its activation condition also holds for $d = 64$. The exact formulas, rather than rounded decimals, define the bounds. The limiting value is $1/3 = 0.333333\dots$



Markers evaluate the proved parameter inequality; joining lines guide the eye.
Table decimals use 100-digit arithmetic and are not interval-certified.

Figure III.97.1: The excess above $1/3$ in two proved upper-bound expressions. The finite expression is evaluated at the parameters stated in Section III.97; the simpler rate is proved for $d \geq 80$. The plotted values are formula evaluations, not estimates of the unknown optimal constant at each finite order.

III.98 An asymptotically sharp quartic cube coefficient

Source 63, delivered Section 5.

The uniform comparison theorem has a direct consequence for the small-perturbation expansion of Gowers cube counts. The quartic coefficient itself was computed in the earlier ProveIt reports [258]. The new conclusion is the exact limiting value of the optimal coefficient when that expansion is expressed in terms of the higher uniformity norm. This is a local refinement of the counting estimates surrounding Lemma 3.10 of Gowers's paper [1].

For a real function g on a finite abelian group, put

$$Q_d(g) = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{\omega \in \{0,1\}^d} g(x + \omega \cdot h) = \|g\|_{U^d}^{2^d}. \quad (\text{III.98.1})$$

All averages are normalized. Write

$$q = 2^d, \quad P_d = \frac{6^d - 2 \cdot 4^d + 2^d}{8}. \quad (\text{III.98.2})$$

The integer P_d counts the four-vertex parallelograms in the discrete d -cube. In particular, $P_2 = 1$ and $P_3 = 12$.

III.98.1 Why odd torsion introduces no quartic correction

The following elementary observation supplies the arithmetic detail needed to work over every odd-order abelian group, including groups with elements of order three.

Lemma III.98.1 (Small binary minors). *Every 3-by-3 matrix with entries in $\{0,1\}$ has determinant of absolute value at most two. Every nonzero 2-by-2 binary determinant has absolute value one.*

Proof. The second assertion follows directly from the determinant formula. For the first, a row with no ones gives determinant zero, and a row with one one reduces the determinant to a 2-by-2 minor. In the remaining case every row is one of

$$110, \quad 101, \quad 011, \quad 111.$$

Repeated rows give zero. If the three distinct rows include 111, coordinate and row permutations reduce to the rows 111, 110, 101, whose determinant has absolute value one. Otherwise the rows are 110, 101, 011, whose determinant has absolute value two. $\square$

Lemma III.98.2 (Four-vertex classification). *Let $V \subseteq \{0,1\}^d$ have four elements. Exactly one of the following holds:*

1. *V is a parallelogram: its elements can be labeled so that $v_1 + v_2 = v_3 + v_4$ in $\mathbb{Z}^d$.*
2. *The four affine forms $x + v \cdot h$, $v \in V$, are jointly uniform on G^4 when $x, h_1, \dots, h_d$ are independent and uniform in any finite abelian group G of odd order.*

Any set of at most three distinct cube vertices gives jointly uniform forms on every finite abelian group.

Proof. Reflect coordinates to move one selected vertex to the origin. This operation preserves the assertions: a reflection in coordinate i replaces the variables by $x' = x + h_i$, $h'_i = -h_i$, an invertible change of uniform variables. The remaining vertex vectors now have entries in $\{0, 1\}$.

Two distinct nonzero binary vectors are linearly independent over $\mathbb{R}$. Otherwise one would be a scalar multiple of the other; a coordinate equal to one forces the scalar to be one and the vectors to coincide. Their rank-two matrix therefore has a nonzero binary 2-by-2 minor, which is a unit by Lemma III.98.1. This proves the assertion for three vertices, and the smaller cases follow in the same way.

For four vertices, consider the rank of the three nonzero vectors. If it is three, a 3-by-3 coordinate minor has determinant ± 1 or ± 2 . Such a matrix acts invertibly on G^3 when $|G|$ is odd: its adjugate supplies an inverse after division by its determinant, and multiplication by two is an automorphism of G . Together with the base variable x , this proves joint uniformity of the four forms.

If the rank is two, project the vertex set onto two coordinates for which the plane has rank two. The projection is injective on the plane, so the four projected vertices are all of $\{0, 1\}^2$. Each remaining coordinate is an affine function on this square taking only the values zero and one. Such a function is constant, one of the two coordinate functions, or its complement: its values at the four corners obey the parallelogram identity. Consequently the original four vertices obey that identity as well. This also shows that a parallelogram cannot fall into the rank-three case. $\square$

Proposition III.98.3 (Exact quartic expansion with a controlled remainder). *Let $d \geq 2$, let G be a finite abelian group of odd order, and let $f : G \rightarrow \mathbb{R}$ satisfy $\mathbb{E}f = 0$. For every $\delta \geq 0$, with $u = \|f\|_{U^d}$,*

$$Q_d(\delta + f) = \delta^q + P_d \delta^{q-4} \|f\|_{U^2}^4 + \mathcal{R}_d(\delta, f), \quad (\text{III.98.3})$$

where

$$|\mathcal{R}_d(\delta, f)| \leq \sum_{j=5}^q \binom{q}{j} \delta^{q-j} u^j. \quad (\text{III.98.4})$$

An empty sum is zero, and a zeroth power is one, including when $\delta = 0$.

Proof. Expand the product in (III.98.1), choosing the vertices at which the factor f is used. The empty choice contributes δ^q . Every choice of one, two, or three vertices contributes zero by joint uniformity and $\mathbb{E}f = 0$.

For a four-vertex choice of affine rank three, the contribution is again zero by Lemma III.98.2. For a parallelogram, the three free vertex forms are jointly uniform by the same lemma, and the fourth is determined by their parallelogram relation. Its contribution is therefore

$$\mathbb{E}_{a,b,c} f(a)f(b)f(c)f(a+b-c) = \|f\|_{U^2}^4.$$

To count these choices, there are 6^d ordered quadruples of binary vertices satisfying $v_1 + v_2 = v_3 + v_4$: in each coordinate there are six allowed bit patterns. The trivial pairings $(v_1, v_2) = (v_3, v_4)$ and $(v_1, v_2) = (v_4, v_3)$ account for $2 \cdot 4^d - 2^d$ quadruples, with their overlap counted once. Every repeated-entry solution is one of these: if a repeated entry lies on both sides it cancels, while $2v_1 = v_3 + v_4$ forces $v_1 = v_3 = v_4$ coordinatewise. Every remaining unordered parallelogram is counted eight times, by interchanging its two diagonals and ordering each diagonal. This gives (III.98.2) and proves the quartic term.

Finally, for a choice of $j \geq 5$ vertices, fill those positions with f and the other cube positions with the constant function one. The mixed Gowers–Cauchy–Schwarz inequality bounds the absolute

value of the resulting average by u^j . There are $\binom{q}{j}$ such choices. Summing their bounds proves (III.98.4). $\square$

The centered expansion and its quartic coefficient are credited to the earlier reports [258]; the short proof above makes the torsion hypotheses explicit. The remainder can have either sign. In particular, (III.98.3) does not imply the lower bound obtained by deleting it. The earlier reports give counterexamples to that deletion already at $d = 3$.

III.98.2 The optimal coefficient in uniformity-norm units

Definition III.98.4. Let A_d be the infimum of constants $A \geq 0$ for which

$$Q_d(\delta + f) \leq \delta^q + A\delta^{q-4}u^4 + \sum_{j=5}^q \binom{q}{j} \delta^{q-j}u^j, \quad u = \|f\|_{U^d}, \quad (\text{III.98.5})$$

holds for every finite abelian group of odd order, every real centered function f , and every $\delta \geq 0$.

Theorem III.98.5 (Exact reduction and asymptotically sharp coefficient). *For every $d \geq 2$,*

$$A_d = P_d c_d^{\text{odd}}. \quad (\text{III.98.6})$$

Consequently,

$$\frac{A_d}{P_d} \rightarrow \frac{1}{3}, \quad A_d \sim \frac{6^d}{24} \quad (d \rightarrow \infty). \quad (\text{III.98.7})$$

For every $d \geq 80$,

$$\frac{P_d}{3} \leq A_d \leq P_d \left(\frac{1}{3} + 10 \cdot 2^{-\lfloor d/8 \rfloor} \right). \quad (\text{III.98.8})$$

The same coefficient is necessary even when sharpness is tested only on perturbations $\delta + tf$ taking values in $[0, 1]$, with a fixed $0 < \delta < 1$ and sufficiently small $t > 0$.

Proof. The definition of c_d^{odd} gives $\|f\|_{U^2}^4 \leq c_d^{\text{odd}}u^4$. Substitution into Proposition III.98.3 proves (III.98.5) with $A = P_d c_d^{\text{odd}}$.

Conversely, suppose a constant A satisfies (III.98.5). Fix an admissible nonzero f , fix $0 < \delta < 1$, and replace f by tf . The finite expansion gives, as $t \downarrow 0$,

$$Q_d(\delta + tf) = \delta^q + P_d \delta^{q-4} t^4 \|f\|_{U^2}^4 + O_{d,\delta,f}(t^5).$$

The assumed upper bound has fourth-order term $A\delta^{q-4}t^4 \|f\|_{U^d}^4$ and a remainder of order t^5 . Subtract δ^q , divide by $\delta^{q-4}t^4$, and let $t \downarrow 0$. This yields

$$P_d \|f\|_{U^2}^4 \leq A \|f\|_{U^d}^4.$$

Taking the supremum over the group and f proves the reverse inequality in (III.98.6). Since the group is finite, f is bounded, and sufficiently small t makes $0 \leq \delta + tf \leq 1$ pointwise. Thus the restricted sharpness assertion follows from the same argument. It concerns real-valued perturbations, without asserting indicator-valued sharpness.

Theorem III.97.1 now proves (III.98.7) and (III.98.8); here $P_d \sim 6^d/8$ by (III.98.2). $\square$

[write] **7 October 2026; sources 60–69.** In the report’s terms this is the best odd-order leading quartic coefficient of item III.1, now asymptotic to $P_d/3$ with $P_d = (6^d - 2 \cdot 4^d + 2^d)/8$ (source 04’s P_d , source 55’s C_d); at $d = 4$ it is $P_4 c_4^{\text{odd}}$ with the bracket above.

For $\delta > 0$, writing $\varepsilon = u/\delta$ displays the small-uniformity scale explicitly:

$$\frac{Q_d(\delta + f)}{\delta^q} \leq 1 + A_d \varepsilon^4 + \sum_{j=5}^q \binom{q}{j} \varepsilon^j. \quad (\text{III.98.9})$$

For each fixed d , the sum after the quartic term is $O_d(\varepsilon^5)$ as $\varepsilon \rightarrow 0$. The asymptotic result (III.98.7) determines how the optimal fourth-order coefficient subsequently behaves as d grows. A simultaneous limit in d and ε must retain the explicit binomial remainder. This refinement does not by itself change a final quantitative bound in Szemerédi’s theorem; such a conclusion would require carrying the improved local estimate through the later inverse and density-increment arguments.

III.99 The limiting real Gowers norm constant

Source 60, delivered Section 2.

The higher-dimensional comparison between real centered Gowers norms has two distinct sources of difficulty: the function may change as the order of the norm increases, and its ambient group may change as well. Taking the limit on each fixed finite group does not justify interchanging that limit with the supremum over groups. We resolve this interchange by controlling the Fourier fourth-power mass and then constructing a compact-group limit.

The resulting theorem identifies the asymptotic comparison constant with a bounded Fourier extremal problem. It also narrows the interval $[1/3, 1/2]$ obtained in source 29 of the repository to $[1/3, 0.3464533951 \dots]$. When the least prime divisor of the group order tends to infinity together with the Gowers norm order, the asymptotic constant is exactly $1/3$. The latter statement uses the published Riesz–Sobolev inequality of Christ and Iliopoulou; the former bound and the compactness reduction have the self-contained proofs below, apart from standard Fourier analysis and Pontryagin duality.

We use the normalized Gowers norms introduced in [1]. Questions concerning their extremizers and nearly extremal functions belong to the broader setting of [7]. The use of varying compact groups as limit objects has substantial precedent, including [241]; we do not claim a new general framework for limits of functions on groups. The specific new input here is the Fourier tail estimate (III.99.11), which allows the starting functions to be unbounded in the supremum norm while their increasingly high Gowers norms remain bounded.

III.99.1 Definitions and statements

All groups in this section are abelian. Finite groups carry uniform probability measure; compact groups carry Haar probability measure. For a real function f define

$$Q_d(f) = \mathbb{E}_{x, h_1, \dots, h_d} \prod_{\omega \in \{0,1\}^d} f(x + \omega \cdot h), \quad \|f\|_{U^d} = Q_d(f)^{1/2^d}. \quad (\text{III.99.1})$$

The standard complex definition inserts conjugation at the odd vertices; for real functions it reduces to (III.99.1). We use the usual triangle inequality, translation invariance, recursion, and monotonicity

of these norms. In particular, for $d \geq 2$, convolution with a nonnegative kernel of integral one contracts U^d .

The normalized Fourier coefficient at a character χ is $\hat{f}(\chi) = \mathbb{E}_x f(x) \overline{\chi(x)}$. The basic identity is

$$\|f\|_{U^2}^4 = \sum_{\chi} |\hat{f}(\chi)|^4. \quad (\text{III.99.2})$$

The sum on the dual group uses counting measure. A function is *centered* if its integral is zero.

For $d \geq 2$ set

$$c_d^{\text{odd}} = \sup_{\substack{G \text{ finite, } |G| \text{ odd} \\ f: G \rightarrow \mathbb{R}, \mathbb{E}f=0, f \neq 0}} \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4}, \quad (\text{III.99.3})$$

$$\beta_{\text{odd}} = \sup_{\substack{G \text{ finite, } |G| \text{ odd} \\ f: G \rightarrow \mathbb{R}, \mathbb{E}f=0, f \neq 0}} \frac{\|f\|_{U^2}^4}{\|f\|_{\infty}^4}. \quad (\text{III.99.4})$$

Both quotients are invariant under nonzero scalar multiplication. Thus the second supremum can equivalently be taken over centered real functions with $|f| \leq 1$.

Theorem III.99.1 (Asymptotic reduction and improved bracket). *The sequence c_d^{odd} is nonincreasing and*

$$\boxed{\lim_{d \rightarrow \infty} c_d^{\text{odd}} = \beta_{\text{odd}}.} \quad (\text{III.99.5})$$

Moreover,

$$\boxed{\frac{1}{3} \leq \beta_{\text{odd}} \leq B := \frac{1}{2} - \frac{8}{\pi^2} + \frac{64}{\pi^4} = 0.3464533951610953\dots} \quad (\text{III.99.6})$$

[write] **7 October 2026; sources 60–69.** *Superseded by source 63* (Section III.94): its Theorem III.96.2 gives $\|f\|_{U^2}^4 \leq \|f\|_{\infty}^4 / 3$ for every real centered f on every finite group of odd order, so $\beta_{\text{odd}} = 1/3$ and the limit in (III.99.5) is $1/3$; source 63 also gives the rate. This theorem is correct as printed; only the bracket (III.99.6) is no longer the best statement.

The exact identity $\beta_{\text{odd}} = 1/3$ is not proved here. Equation (III.99.5) makes it equivalent to the exact asymptotic question for the Gowers norm comparison. The upper bound is qualitative as a statement about c_d^{odd} : for every $\varepsilon > 0$ there exists d_0 such that $c_d^{\text{odd}} \leq B + \varepsilon$ for $d \geq d_0$, but the argument below supplies no effective value of d_0 . It therefore makes no new claim about the fourth-order coefficient 4795/832 in source 45.

For a nontrivial finite group G , write $p_*(G)$ for the least prime divisor of $|G|$.

Theorem III.99.2 (Exact limit as small prime torsion disappears). *Suppose $d_n \rightarrow \infty$, $p_*(G_n) \rightarrow \infty$, and $f_n : G_n \rightarrow \mathbb{R}$ are nonzero centered functions. Then*

$$\limsup_{n \rightarrow \infty} \frac{\|f_n\|_{U^2}^4}{\|f_n\|_{U^{d_n}}^4} \leq \frac{1}{3}. \quad (\text{III.99.7})$$

The constant is sharp. More explicitly, if

$$C(d, P) = \sup_{\substack{G \text{ finite, } p_*(G) \geq P \\ f: G \rightarrow \mathbb{R}, \mathbb{E}f=0, f \neq 0}} \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4}, \quad (\text{III.99.8})$$

then

$$\boxed{\lim_{\substack{d \rightarrow \infty \\ P \rightarrow \infty}} C(d, P) = \frac{1}{3}} \quad (\text{III.99.9})$$

in the joint sense: for every $\varepsilon > 0$, some d_0, P_0 work for all $d \geq d_0$ and $P \geq P_0$.

[write] **7 October 2026; sources 60–69.** Contained in source 63's Theorem III.97.1: if $p_*(G) \geq 3$ then $|G|$ is odd, so $C(d, P) \leq c_d^{\text{odd}} \leq 1/3 + 10 \cdot 2^{-\lfloor d/8 \rfloor}$ for $d \geq 80$ and every $P \geq 3$, with no condition on P (checked at the write). The Riesz–Sobolev route remains a second proof of the joint limit.

Theorem III.99.2 concerns a simultaneous limit. It does not assert $c_d^{\text{odd}} \leq 1/3$ at a fixed order, nor a bound $1/3$ uniform in d for a fixed finite prime modulus.

III.99.2 Positive moments prevent escape of Fourier mass

For a real function f on a finite group, its autocorrelation

$$r_f(h) = \mathbb{E}_x f(x) f(x+h) \quad (\text{III.99.10})$$

has Fourier coefficients $|\widehat{f}(\chi)|^2 \geq 0$. Consequently every surviving term in the Fourier expansion of an even autocorrelation moment has a nonnegative coefficient. Source 29 uses balanced words in this expansion to obtain the central-binomial moment bound. Here we retain words using *distinct* conjugate pairs; this controls the tail of the spectrum rather than only its total mass.

Lemma III.99.3 (Uniform Fourier tail estimate). *Let $f : G \rightarrow \mathbb{R}$ be centered, let $|G|$ be odd, and suppose $\|f\|_{U^D} \leq 1$, where $D \geq 2$. Choose one character χ_j from each nontrivial inverse pair, ordering the pairs so that*

$$x_j := |\widehat{f}(\chi_j)|^4 \quad \text{satisfies} \quad x_1 \geq x_2 \geq \cdots.$$

Append zero values if the list is finite. For every integer $L \geq 0$ and every integer $1 \leq m \leq 2^{D-2}$,

$$\begin{aligned} 2 \sum_{j>L} x_j &\leq 2 \left(\frac{m!}{(2m)!} \right)^{1/m} + \frac{m-1}{L+1} \\ &\leq \frac{2}{m+1} + \frac{m-1}{L+1}. \end{aligned} \quad (\text{III.99.11})$$

Proof. The group has no nontrivial self-inverse characters, and centering removes the trivial one. Thus

$$2 \sum_j x_j = \|f\|_{U^2}^4 \leq 1. \quad (\text{III.99.12})$$

Put $M = 2^{D-2}$. The derivative recursion and $\|g\|_{U^{D-1}} \geq |\mathbb{E}g|$ imply

$$\mathbb{E}_h |r_f(h)|^{2M} \leq \|f\|_{U^D}^{2D} \leq 1.$$

Monotonicity of the L^p norms on a probability space gives $\mathbb{E}|r_f|^{2m} \leq 1$ for every $m \leq M$.

In the Fourier expansion of this last moment, retain those words that choose m distinct pairs from indices greater than L and use each character and its inverse exactly once. If e_m denotes the elementary symmetric polynomial, their contribution is $(2m)!e_m(x_{L+1}, x_{L+2}, \dots)$. Therefore

$$1 \geq \frac{(2m)!}{m!} \mathcal{D}_m, \quad \mathcal{D}_m = \sum_{\substack{j_1, \dots, j_m > L \\ j_1, \dots, j_m \text{ distinct}}} x_{j_1} \cdots x_{j_m}. \quad (\text{III.99.13})$$

Extra character relations create further nonnegative terms, so they do not invalidate this lower bound.

Write

$$T = \sum_{j>L} x_j, \quad a = \max_{j>L} x_j \leq \frac{1}{2(L+1)},$$

where the last bound follows from the ordering and (III.99.12). If $T > (m-1)a$, successive choices of distinct indices leave total mass at least $T - ja$ after j choices. Induction gives

$$\mathcal{D}_m \geq \prod_{j=0}^{m-1} (T - ja) \geq (T - (m-1)a)^m.$$

Combining this with (III.99.13) yields

$$T \leq \left(\frac{m!}{(2m)!} \right)^{1/m} + (m-1)a.$$

If $T \leq (m-1)a$, the same conclusion is immediate. Multiplication by two proves the first inequality in (III.99.11). For the second, use $(2m)!/m! = \prod_{j=m+1}^{2m} j \geq (m+1)^m$. $\square$

For a sequence $D_n \rightarrow \infty$, the right side of (III.99.11) tends to zero uniformly in all sufficiently large n by first choosing m large and then L large. This is the precise tightness statement needed below. Choosing m of order $\sqrt{L}$ also gives an explicit tail of order $L^{-1/2}$ whenever D is large enough to permit that choice of m .

III.99.3 The limit of the norms on a fixed compact group

The next lemma is used only for real trigonometric polynomials. Its proof keeps track of signs; restricting a cube average to a set of directions without doing so would not be legitimate.

Lemma III.99.4 (Fixed trigonometric polynomials). *If g is a real trigonometric polynomial on a compact abelian group K , then*

$$\lim_{d \rightarrow \infty} \|g\|_{U^d} = \|g\|_{\infty}. \quad (\text{III.99.14})$$

Proof. Let $M = \|g\|_{\infty}$. The zero case is immediate. Fix $0 < \varepsilon < M/4$, and let

$$E = \{x : |g(x)| \geq M - \varepsilon\}, \quad u = \mu(E) > 0.$$

The positivity of u follows from continuity and the full support of Haar measure. Write $g = \sum_{j=1}^r a_j \chi_j$ and put $A = \sum_j |a_j|$. Choose a fixed $0 < \eta < 1/8$ with $2\pi A\eta < \varepsilon$. Define

$$B_d = \{h : \|\chi_j(h)\|_{\mathbb{R}/\mathbb{Z}} < \eta/d \text{ for } 1 \leq j \leq r\},$$

where a unit complex number is identified with its phase modulo one. If $h_1, \dots, h_{d-1} \in B_d$, then for every $\omega \in \{0, 1\}^{d-1}$ and every x ,

$$|g(x + \omega \cdot h) - g(x)| < \varepsilon. \quad (\text{III.99.15})$$

Indeed the displacement of each character along a sum is bounded by the sum of its displacements along the individual increments.

We shall use the elementary estimate

$$\mu(B_d) \geq \left(\frac{\eta}{3d}\right)^r. \quad (\text{III.99.16})$$

To see this, divide each coordinate of the image of K in $(\mathbb{R}/\mathbb{Z})^r$ into $N = \lceil 2d/\eta \rceil$ intervals. One inverse image of a box has measure at least N^{-r} . Subtracting any point in that inverse image places its translate inside B_d . Thus $\mu(B_d) \geq N^{-r} \geq (\eta/(3d))^r$.

For $k = 2^{d-1}$ and a tuple of directions h put

$$R(h) = \mathbb{E}_x \prod_{\omega \in \{0,1\}^{d-1}} g(x + \omega \cdot h).$$

If the directions lie in B_d and $|g(x)| > \varepsilon$, all factors have the same sign by (III.99.15). Since k is even for $d \geq 2$, their product is nonnegative. On E its value is at least $(M - 2\varepsilon)^k$. Where $|g(x)| \leq \varepsilon$, its negative part is at most $(2\varepsilon)^k$. Consequently

$$R(h) \geq u(M - 2\varepsilon)^k - (2\varepsilon)^k \geq \frac{u}{2}(M - 2\varepsilon)^k$$

for all sufficiently large d and all $h \in B_d^{d-1}$.

Averaging the final direction in the cube gives the exact identity

$$Q_d(g) = \mathbb{E}_{h_1, \dots, h_{d-1}} R(h)^2. \quad (\text{III.99.17})$$

Indeed the two parallel faces have base points x and $x + h_d$, which become independent uniform points after averaging x, h_d . We may therefore restrict this nonnegative square average to B_d^{d-1} . By (III.99.16),

$$\|g\|_{U^d} \geq (M - 2\varepsilon) \left[\left(\frac{\eta}{3d}\right)^{r(d-1)} \left(\frac{u}{2}\right)^2 \right]^{1/2^d}.$$

The factor in brackets raised to $1/2^d$ tends to one. The opposite bound $\|g\|_{U^d} \leq M$ is immediate from the definition. Letting ε tend to zero proves the claim. $\square$

III.99.4 A compact limit from high-order norm control

We now prove the main compactness statement. The limit group can have elements of order two; for example, the circle does. The property preserved from odd-order finite groups is instead that its *discrete dual* has no nonzero element of order two. Keeping this distinction explicit is essential for the finite approximation later.

Theorem III.99.5 (Compactness with preservation of the second norm). *Let G_n be finite odd-order abelian groups, let $D_n \rightarrow \infty$, and let $f_n : G_n \rightarrow \mathbb{R}$ be centered with $\|f_n\|_{U^{D_n}} \leq 1$. Pass to a subsequence on which $\|f_n\|_{U^2}^4$ converges. There exist a compact metrizable abelian group K , whose discrete dual has no nonzero element of order two, and a measurable centered real function f with $|f| \leq 1$ almost everywhere, such that*

$$\|f\|_{U^2}^4 = \lim_{n \rightarrow \infty} \|f_n\|_{U^2}^4. \quad (\text{III.99.18})$$

If additionally $p_(G_n) \rightarrow \infty$, K can be taken connected.*

Proof. If the orders $|G_n|$ remain bounded, take a further subsequence with a fixed finite group. Every Fourier coefficient has modulus at most one, so there is a coefficientwise limit f . For every fixed d , eventually $D_n \geq d$, whence $\|f_n\|_{U^d} \leq 1$ and therefore $\|f\|_{U^d} \leq 1$. Lemma III.99.4 gives $\|f\|_\infty \leq 1$; all other claims follow by finite Fourier continuity. We may henceforth assume $|G_n| \rightarrow \infty$.

Step 1: stabilize the character relations. For each n , choose representatives $\chi_{n,j}$ of the nontrivial inverse pairs in decreasing order of $|\widehat{f}_n(\chi_{n,j})|$. Pad the infinite list by trivial characters after all genuine pairs have been listed. Every fixed initial segment consists of distinct nontrivial inverse pairs for all sufficiently large n .

Let $F = \bigoplus_{j \geq 1} \mathbb{Z}e_j$ be a countable free abelian group, and let

$$\phi_n : F \longrightarrow \widehat{G}_n, \quad \phi_n(e_j) = \chi_{n,j}.$$

The group F is countable. Since every Fourier coefficient has modulus at most one, a diagonal subsequence has the following two properties for every $w \in F$: the truth value of $\phi_n(w) = 0$ is eventually constant, and $\widehat{f}_n(\phi_n(w))$ converges. Let

$$R = \{w \in F : \phi_n(w) = 0 \text{ for all sufficiently large } n\}, \quad \Gamma = F/R.$$

The set R is a subgroup. If $2w \in R$, oddness of $|G_n|$ forces $w \in R$, so Γ has no nonzero element of order two. Set $K = \widehat{\Gamma}$. This is compact and metrizable. For $\gamma = w + R$ define

$$b_\gamma = \lim_{n \rightarrow \infty} \widehat{f}_n(\phi_n(w)). \quad (\text{III.99.19})$$

This is independent of the representative. Moreover $b_0 = 0$ and $b_{-\gamma} = \overline{b_\gamma}$.

Step 2: positive polynomial kernels produce bounded limits. Write $\gamma_j = e_j + R$, viewing γ_j also as a character on K . For fixed positive integers J, t , define

$$P_{J,t}(x) = \prod_{j=1}^J \left| \frac{1 + \gamma_j(x)}{2} \right|^{2t}, \quad K_{J,t} = \frac{P_{J,t}}{\int_K P_{J,t}}. \quad (\text{III.99.20})$$

The denominator is positive: $P_{J,t}$ is continuous and equals one at the identity. The kernel is real, nonnegative, has integral one, and has finite Fourier support.

Substituting $\chi_{n,j}$ for γ_j in the same expression gives a nonnegative kernel on G_n , normalized to have mean one, denoted $K_{n,J,t}$. Only finitely many words in F occur in the polynomial. Their relations stabilize. Hence the normalizing constants are eventually the same as the constant in (III.99.20), and all identifications of frequencies in these polynomials eventually agree.

The functions $g_n = f_n * K_{n,J,t}$ satisfy $\|g_n\|_{U^d} \leq 1$ for every fixed d and all large n . This follows from the norm triangle inequality and translation invariance, because the kernel is a probability kernel and $D_n \geq d$. Their Fourier support is contained in a fixed finite set of formal words. The coefficients converge to those of the real trigonometric polynomial $g_{J,t}$ on K defined by

$$\widehat{g}_{J,t}(\gamma) = b_\gamma \widehat{K}_{J,t}(\gamma). \quad (\text{III.99.21})$$

For each fixed d , expansion of a d -cube average in these finitely many Fourier coefficients involves finitely many character relations. All such relations have stabilized. Fourier orthogonality therefore gives $Q_d(g_n) \rightarrow Q_d(g_{J,t})$ and $\|g_{J,t}\|_{U^d} \leq 1$. This holds for every fixed d . Lemma III.99.4 now yields

$$\|g_{J,t}\|_\infty \leq 1. \quad (\text{III.99.22})$$

Step 3: remove the smoothing. For fixed J , write the numerator of (III.99.20) as $v_J(x)^t$, where $0 \leq v_J \leq 1$ and $v_J = 1$ exactly on

$$H_J = \{x : \gamma_j(x) = 1 \text{ for } 1 \leq j \leq J\}.$$

Outside any fixed neighborhood of H_J , the continuous function v_J is at most some $c < 1$. On a smaller neighborhood of the identity of positive measure, it exceeds a number c' with $c < c' < 1$. It follows that the normalized kernel mass outside the first neighborhood is at most a constant times $(c/c')^t$. Thus, whenever $\gamma \in \langle \gamma_1, \dots, \gamma_J \rangle$,

$$\widehat{K}_{J,t}(\gamma) \longrightarrow 1 \quad (t \rightarrow \infty). \quad (\text{III.99.23})$$

Every element of Γ belongs to one of these finitely generated subgroups. A diagonal choice of J and t produces a sequence of kernels for which (III.99.23) holds at every $\gamma \in \Gamma$. The associated polynomials $g_{J,t}$ are centered and real, and satisfy (III.99.22). The unit ball of $L^\infty(K)$ is weak-star compact. Since K is compact metrizable, $L^1(K)$ is separable, so we may take a weak-star convergent subsequence. Its limit f is real centered, has $|f| \leq 1$ almost everywhere, and has Fourier coefficients b_γ .

Step 4: recover all the Fourier fourth-power mass. The elements γ_j are distinct, nonzero, and no two are inverses: these properties hold for each fixed initial list in all sufficiently large finite groups, so they hold after stabilization. If $b_\gamma \neq 0$, choose $\varepsilon > 0$ smaller than its modulus. The corresponding coefficient of f_n has modulus at least ε for all large n . Its inverse pair must lie among at most $1/(2\varepsilon^4)$ of the ordered pairs, because $\|f_n\|_{U^2}^4 \leq 1$. Stabilization forces $\gamma = \gamma_j$ or $\gamma = -\gamma_j$ for some j in this fixed finite list. Indeed if none of these finitely many equalities held in Γ , all would eventually fail in the finite groups, a contradiction. Thus all nonzero coefficients b_γ occur in the listed pairs.

For each fixed m, L , Lemma III.99.3 applies to all sufficiently large n . First choose m so that $2/(m+1)$ is small, then choose L so that $(m-1)/(L+1)$ is small. The tails of the Fourier fourth-power sums are therefore uniformly small. Passing to the limit in finite initial sums and then removing the tails gives

$$\sum_{\gamma \in \Gamma} |b_\gamma|^4 = \lim_{n \rightarrow \infty} 2 \sum_j |\widehat{f}_n(\chi_{n,j})|^4.$$

Using (III.99.2) proves (III.99.18).

Finally, if $p_*(G_n) \rightarrow \infty$ and $mw \in R$ for a fixed positive integer m , then m is coprime to $|G_n|$ for all large n . Consequently $w \in R$. The group Γ is now torsion-free. Pontryagin duality says that $\widehat{\Gamma}$ is connected exactly in this situation, proving the last assertion. $\square$

III.99.5 Finite odd groups approximate the bounded limit problem

Lemma III.99.6 (Finite approximation). *Let K be a compact abelian group whose discrete dual has no nonzero element of order two. If f is centered real and $|f| \leq 1$ almost everywhere, then*

$$\|f\|_{U^2}^4 \leq \beta_{\text{odd}}. \quad (\text{III.99.24})$$

Proof. Choose finitely many Fourier coefficients whose fourth powers sum to at least $\|f\|_{U^2}^4 - \varepsilon$, and let H be the subgroup of $\widehat{K}$ that they generate. Conditional expectation over the annihilator of H

yields a function on $\widehat{H}$ preserving these coefficients, the mean, realness, and the bound one. The structure theorem for finitely generated abelian groups gives

$$H \cong \mathbb{Z}^r \oplus T, \quad \widehat{H} \cong (\mathbb{R}/\mathbb{Z})^r \times \widehat{T},$$

where T is finite of odd order. Oddness follows from the absence of order-two elements in H .

Convolution with positive finite Fourier kernels, for example the kernels (III.99.20) formed from a finite generating set of H , gives a continuous real centered function h on $\widehat{H}$ with $|h| \leq 1$ and with the selected Fourier coefficients arbitrarily close to their previous values. We may therefore arrange

$$\|h\|_{U^2}^4 \geq \|f\|_{U^2}^4 - 2\varepsilon.$$

For odd integers $N \rightarrow \infty$, restrict h to the finite subgroups

$$H_N = ((1/N)\mathbb{Z}/\mathbb{Z})^r \times \widehat{T}.$$

Riemann sums converge to Haar integrals for h and its continuous two-cube product. If μ_N is the mean of $h|_{H_N}$, then $\mu_N \rightarrow 0$, and

$$h_N = \frac{h|_{H_N} - \mu_N}{1 + |\mu_N|}$$

is centered real and bounded by one. Its second Gowers norm converges to that of h . Since each H_N has odd order, the definition of β_{odd} gives $\beta_{\text{odd}} \geq \|f\|_{U^2}^4 - 2\varepsilon$. Let ε tend to zero. $\square$

Proof of (III.99.5). Monotonicity of the Gowers norms makes c_d^{odd} nonincreasing. The bound $\|f\|_{U^d} \leq \|f\|_{\infty}$ gives $c_d^{\text{odd}} \geq \beta_{\text{odd}}$ for every d . Let c be the limit of this decreasing sequence.

Choose $D_n \rightarrow \infty$ and functions whose ratios in (III.99.3) tend to c . Rescale them so that $\|f_n\|_{U^{D_n}} = 1$. Then $\|f_n\|_{U^2}^4 \rightarrow c$. Theorem III.99.5 gives a centered real function f bounded by one on a compact group with no dual order-two element, such that $\|f\|_{U^2}^4 = c$. Lemma III.99.6 gives $c \leq \beta_{\text{odd}}$, proving equality. $\square$

III.99.6 The bounded Fourier extremum

The numerical upper bound in Theorem III.99.1 uses a sharp single-coefficient estimate together with the fact that real Fourier coefficients occur in inverse pairs. Centering is essential for the sharp finite-order estimate.

Lemma III.99.7 (A sharp centered coefficient bound). *Let f be real centered on a finite group, with $|f| \leq 1$. If a character χ has odd order $q \geq 3$, then*

$$|\widehat{f}(\chi)| \leq a_q := \frac{1}{q} \cot \frac{\pi}{2q} < \frac{2}{\pi}. \quad (\text{III.99.25})$$

The finite- q constant is attained.

Proof. Conditional expectation on the quotient generated by χ reduces the question to a vector $(v_0, \dots, v_{q-1})$ with $v_j \in [-1, 1]$ and $\sum_j v_j = 0$. Write $q = 2n + 1$. The extreme points of this polytope have n coordinates equal to one, n equal to minus one, and one equal to zero. At most one coordinate can lie strictly inside $[-1, 1]$ at an extreme point, since otherwise two coordinates can be perturbed in opposite directions. The zero sum and the odd number of coordinates then force the remaining coordinate to be zero and the two signs to occur equally often.

Fix an angle θ and maximize

$$\sum_{j=0}^{q-1} v_j \cos(2\pi j/q - \theta).$$

A maximizing extreme point places its positive coordinates at the n largest cosines, its negative coordinates at the n smallest, and its zero at a median. The n largest cosines correspond to an arc of n consecutive roots of unity, and a median root is adjacent to that arc. This follows by ordering the roots by circular distance from θ ; ties can be broken in this way without changing the maximum.

Up to rotation and reflection, its Fourier numerator is therefore

$$1 + 2 \sum_{j=1}^n e^{2\pi i j/q} = i \cot \frac{\pi}{2q}. \quad (\text{III.99.26})$$

Taking θ to be the phase of any candidate Fourier coefficient shows that its modulus is at most the modulus in (III.99.26), divided by q . The vector with zero at 0, positives at $1, \dots, n$, and negatives at the remaining indices attains the bound. Finally $\cot t < 1/t$ for $0 < t < \pi/2$ proves the last inequality. $\square$

Proposition III.99.8 (The Fourier mass cap). *Every centered real function f on a finite odd-order group satisfies*

$$\|f\|_{U^2}^4 \leq \left(\frac{1}{2} - \frac{8}{\pi^2} + \frac{64}{\pi^4} \right) \|f\|_{\infty}^4. \quad (\text{III.99.27})$$

Proof. Normalize $\|f\|_{\infty} = 1$, ignoring the zero function. Choose one representative of each nontrivial inverse pair and set $\lambda_j = |\hat{f}(\chi_j)|^2$. Parseval and centering give

$$\sum_j \lambda_j = \frac{1}{2} \|f\|_2^2 \leq \frac{1}{2}.$$

Lemma III.99.7 gives $\lambda_j \leq A := 4/\pi^2$, with $1/4 < A < 1/2$.

If $\lambda_{\max} \leq 1/4$, then $2 \sum_j \lambda_j^2 \leq \lambda_{\max} \leq 1/4$. Otherwise write $a = \lambda_{\max} \in [1/4, A]$. The remaining mass is at most $1/2 - a$, so

$$2 \sum_j \lambda_j^2 \leq 2 \left[a^2 + (1/2 - a)^2 \right] \leq 2 \left[A^2 + (1/2 - A)^2 \right].$$

The second inequality uses that the expression is increasing for $a \in [1/4, 1/2]$; its value at A exceeds $1/4$, so it covers the first case as well. By (III.99.2), the left side is $\|f\|_{U^2}^4$. Expanding the final expression proves (III.99.27). $\square$

The lower bound $1/3$ admits a bounded finite witness with an error of order q^{-2} . On $G = \mathbb{Z}/q\mathbb{Z}$, where $q = 2n + 1$, set

$$f_q(0) = 0, \quad f_q(j) = 1 \ (1 \leq j \leq n), \quad f_q(j) = -1 \ (n+1 \leq j \leq 2n). \quad (\text{III.99.28})$$

This function is centered and has supremum norm one. Its unnormalized autocorrelation $R(h) = \sum_x f_q(x) f_q(x+h)$ is

$$R(0) = q - 1, \quad R(h) = R(q-h) = q - 4h \quad (1 \leq h \leq n).$$

Indeed for such a nonzero h , exactly two terms contain the zero entry. Among the other $q - 2$ terms there are $2h - 1$ pairs of opposite signs: the two sign intervals meet after translation across their two boundaries, with h crossings at one boundary and $h - 1$ at the boundary containing the zero. Thus $R(h) = (q - 2) - 2(2h - 1) = q - 4h$. It follows that

$$\sum_h R(h)^2 = (q - 1)^2 + 2 \sum_{h=1}^n (q - 4h)^2 = \frac{q^3 - 4q + 3}{3}.$$

Here the sum of squares is obtained directly from $\sum_{h=1}^n h = n(n + 1)/2$ and $\sum_{h=1}^n h^2 = n(n + 1)(2n + 1)/6$. Fourier orthogonality, or the autocorrelation form of U^2 , now gives

$$\|f_q\|_{U^2}^4 = \frac{q^3 - 4q + 3}{3q^3} = \frac{1}{3} - \frac{4}{3q^2} + \frac{1}{q^3} \longrightarrow \frac{1}{3}. \quad (\text{III.99.29})$$

Since $\|f_q\|_{U^d} \leq 1$, the right side is also a lower bound for the ratio against U^d , for every d . This has a smaller finite-size error than the scaled centered-interval model of source 29, which has an error of order q^{-1} . No global optimality of this cyclic witness is needed or asserted. Together with Proposition III.99.8 and (III.99.5), it completes the proof of Theorem III.99.1.

III.99.7 An exact connected-group bound and the joint limit

The next step imports precisely one rearrangement inequality. Theorem 1.2, equation (1.5), of Christ and Iliopoulou [242] states that, for measurable sets A, B, C in a compact connected abelian group, with measures a, b, c satisfying

$$|a - b| \leq c \leq a + b, \quad a + b + c \leq 2,$$

one has

$$\int_C \mathbf{1}_A * \mathbf{1}_B \leq ab - \frac{1}{4}(a + b - c)^2. \quad (\text{III.99.30})$$

Their formulation is an equivalent form of an earlier inequality of Tao, as explained in that paper. We use the inequality as an established result and give the consequence needed here in full.

Lemma III.99.9 (Bounded centered functions on connected groups). *If K is compact connected abelian and f is real centered with $|f| \leq 1$ almost everywhere, then*

$$\|f\|_{U^2}^4 \leq \frac{1}{3}. \quad (\text{III.99.31})$$

The constant is attained on the circle.

Proof. First suppose $f = 2\mathbf{1}_A - 1$, where $\mu(A) = 1/2$, and put $u = \mathbf{1}_A * \mathbf{1}_{-A}$. Then $0 \leq u \leq 1/2$. For any measurable C with $\mu(C) = c$, the hypotheses of (III.99.30) hold with $a = b = 1/2$, so

$$\int_C u \leq \frac{c}{2} - \frac{c^2}{4}. \quad (\text{III.99.32})$$

For $0 \leq t \leq 1/2$, the variational formula for the positive part gives

$$\begin{aligned} \int_K (u - t)_+ &= \sup_C \left(\int_C u - t\mu(C) \right) \\ &\leq \max_{0 \leq c \leq 1} \left((1/2 - t)c - c^2/4 \right) = (1/2 - t)^2. \end{aligned}$$

Integrating this inequality in t yields

$$\int_K u^2 = 2 \int_0^{1/2} \int_K (u-t)_+ dt \leq 2 \int_0^{1/2} (1/2-t)^2 dt = \frac{1}{12}.$$

Since the zero Fourier coefficient of $\mathbf{1}_A$ is $1/2$,

$$\|2\mathbf{1}_A - 1\|_{U^2}^4 = 16 \left(\int_K u^2 - \frac{1}{16} \right) \leq \frac{1}{3}.$$

For general f , write $g = (f + 1)/2$ and form the measurable subgraph

$$A = \{(x, t) \in K \times (\mathbb{R}/\mathbb{Z}) : 0 \leq t < g(x)\},$$

using representatives $t \in [0, 1)$. This set has Haar measure $1/2$, and the product group is compact connected. Averaging $2\mathbf{1}_A - 1$ over translations in the circle subgroup gives the function $(x, t) \mapsto f(x)$. Such averaging contracts U^2 . The sign-function case just proved therefore gives (III.99.31). For a half-circle interval, its autocorrelation is the triangular function and the inequalities above are equalities; equivalently, the centered circle square wave has Fourier fourth-power sum $1/3$. $\square$

Proof of Theorem III.99.2. Normalize $\|f_n\|_{U^{d_n}} = 1$ and extract a subsequence on which $\|f_n\|_{U^2}^4$ converges to the limsup in question. Theorem III.99.5 gives a compact-group limit f bounded by one, preserving this fourth power. Its dual is torsion-free because $p_*(G_n) \rightarrow \infty$, so its group is connected. Lemma III.99.9 proves the upper bound $1/3$.

The cyclic construction (III.99.29) may be taken along odd primes $2n+1 \rightarrow \infty$. Since $\|f_n\|_{U^{d_n}} \leq \|f_n\|_\infty$, its ratios against the increasing-order Gowers norms have liminf at least $1/3$, for any accompanying sequence $d_n \rightarrow \infty$. The upper bound already proved then forces convergence to $1/3$. For every fixed d, P , the same interval examples show $C(d, P) \geq 1/3$. If the joint upper limit in (III.99.9) exceeded $1/3$, one could choose near-maximizers with $d_n \geq n$ and $p_*(G_n) \geq n$ violating (III.99.7). This proves the joint formulation. $\square$

III.99.8 A bounded-exponent variant

The compactness argument also records persistent torsion rather than discarding it. Fix an odd integer $q \geq 3$. Define $c_d(q)$ and $\beta(q)$ as in (III.99.3)–(III.99.4), restricting to finite groups whose exponent divides q .

Corollary III.99.10 (Fixed exponent). *For each fixed odd $q \geq 3$,*

$$\lim_{d \rightarrow \infty} c_d(q) = \beta(q) \leq \frac{1}{2} - 2A_q + 4A_q^2, \quad A_q = \left(\frac{1}{q} \cot \frac{\pi}{2q} \right)^2. \quad (\text{III.99.33})$$

Proof. In the marked-dual construction every relation $qe_j = 0$ is preserved, so the limit dual also has exponent dividing q . Its finitely generated subgroups are finite with exponent dividing q . Lemma III.99.6 therefore returns to finite groups in the same class, with no torus approximation needed. The proof of (III.99.5) gives the asserted equality.

Every nontrivial character has odd order at most q . The quantity $a_r = r^{-1} \cot(\pi/(2r))$ is increasing in r : after putting $t = \pi/(2r)$, this follows from

$$\frac{d}{dt}(t \cot t) = \frac{\sin t \cos t - t}{\sin^2 t} < 0.$$

Thus Lemma III.99.7 bounds every pair mass by A_q . Since $1/4 < A_q < 1/2$, the same mass-cap calculation as in Proposition III.99.8 proves (III.99.33). $\square$

For example, the upper bounds for $q = 3, 5, 7$ are respectively

$$\frac{5}{18} = 0.277777\dots, \quad 0.3164458247200069\dots, \quad 0.3303703837266131\dots$$

These are bounds for the limits within fixed exponent classes and are not asserted sharp. Their being below $1/3$ is consistent with Theorem III.99.2: the latter allows the characteristic to tend to infinity, whereas these classes keep it fixed.

The compactness proof separates Fourier tail control, algebraic relations, and boundedness of the limit. The resulting extremal and quantitative questions are developed in Section V.226.

III.100 Source 64's chapter: binomial separation at Gowers's Section 4 [write]

Added in the eleventh write, 7 October 2026.

The next seven sections are source 64's Sections 2–8, on the quantitative form of Gowers's Section 4 conjectures: how small the progression count of a set can be under one degree of uniformity less than the counting theorem needs. It is the report's first chapter on Section 4. *It uses the colouring material* of source 10's Section III.30, source 14's Section III.36 and source 22's chapter (Section III.60); source 64's introduction is Section F.69.

What is proved. $D = \{0, 1, 4, 13\}$ is $(1, 5, 10)$ -separating modulo 80 (Proposition III.101.4), so digit lifting gives sets with no nontrivial solution of the six-binomial equation at reciprocal growth exponent $\log_4 80 = 3.1609640$, and $3 \leq \beta_6 \leq \log_4 80$ (Theorem III.101.5, Corollary III.101.7), against the exponent 5 of Deng–Tidor–Zhao's general construction; no non-singleton cyclic alphabet is perfect for $(1, 5, 10)$ (Theorem III.102.1, Corollary III.102.2). A symmetric-progression-free r -colouring of an interval of length N transfers to a $9r$ -colouring of the circle whose symmetric k -progression probability is exactly $1/(9(k-1)N)$ (Theorem III.103.2), and with mirror-separating labels to the torus with an exact progression density (Theorem III.104.1); with the prime-indexed colourings of Shi and Dong, credited as a 2026 preprint input, $\rho_4(\alpha) \ll \alpha^5$ and $\rho_6(\alpha) \ll \alpha^{5+9/\log_4 80} = \alpha^{7.847}$ with no ε loss (Theorem III.105.1), realized by deterministic sets along primes (Theorem III.106.1, Corollary III.106.3). The disproofs of the conjectures themselves are Gowers's (length four) and Shi and Dong's, not claimed here. *The dimension barrier* (Theorem III.107.2: a degree- $(k-1)$ colouring avoiding symmetric k -progressions has $n \leq tk^2/4$, sharp for scalar colours) *is source 10's exact threshold* (Theorem III.30.3) at $D_\ell = k-1$, where $s_\ell = k/2$ and $B = tk^2/4$, which source 10 proves with sharpness for every t and source 14 reproves (Theorem III.36.1); source 64 does not cite either, and says that it makes no first-discovery claim for polynomial-method consequences of this kind (note after the theorem).

Relation to the formal statements (checked at HEAD fb9602e55). `conjecture_4_1` and `conjecture_4_2` (Section04.lean, lines 30 and 39) are proposition-valued definitions, not asserted, with progression length $k+2$ in the latter, as source 64 says; nothing in this chapter is formalized.

III.101 Weighted separation and the six-variable construction

Source 64, delivered Section 2.

III.101.1 Trivial solutions really are mirror solutions at length six

Definition III.101.1. For positive integer weights $w = (w_1, \dots, w_h)$, a subset D of an abelian group is *w-separating* if the map

$$D^h \longrightarrow G, \quad (d_1, \dots, d_h) \longmapsto \sum_{j=1}^h w_j d_j$$

is injective. Equality of ordered tuples, rather than equality of multisets, is required.

The six-binomial equation is equivalently

$$x_0 + 5x_4 + 10x_2 = x_5 + 5x_1 + 10x_3. \quad (\text{III.101.1})$$

Thus $(1, 5, 10)$ -separation is precisely the desired mirror conclusion. There is no unmentioned exceptional family of trivial solutions in this case.

Lemma III.101.2. *In the usual definition of a trivial solution to an invariant integer equation, namely that the coefficients attached to each distinct value sum to zero, the trivial solutions of (F.69.5) are exactly the mirror solutions. The analogous assertion holds for $x_0 - 3x_1 + 3x_2 - x_3 = 0$.*

Proof. The subset sums of $\{1, 5, 10\}$ are distinct: a relation $\varepsilon_1 + 5\varepsilon_2 + 10\varepsilon_3 = 0$ with $\varepsilon_j \in \{-1, 0, 1\}$ first gives $\varepsilon_3 = 0$, because $10 > 1 + 5$, and then $\varepsilon_2 = \varepsilon_1 = 0$. For a fixed value among the six coordinates, cancellation of its positive and negative coefficients therefore requires each weight to occur on both sides or on neither. Applying this to every value gives the three mirror equalities. Conversely those equalities cancel all coefficients. Replace the weight list by $(1, 3)$ for the four-variable statement. $\square$

Remark III.101.3. The proof uses a property of these particular weight lists. At other even lengths the positive half-binomial coefficients need not have distinct subset sums. One must not silently replace “trivial” by “mirror” for every even length; [37, Remark 3.5] explicitly identifies a failure at length fourteen.

III.101.2 The finite certificate

Proposition III.101.4 (A four-symbol seed modulo eighty). *The set*

$$D = \{0, 1, 4, 13\} \subseteq \mathbb{Z}/80\mathbb{Z}$$

is $(1, 5, 10)$ -separating.

Proof. Suppose $a + 5b + 10c \equiv a' + 5b' + 10c' \pmod{80}$ with all six entries in D . The residues of D modulo 5 are distinct, so reduction modulo 5 gives $a = a'$ as chosen integer representatives. Dividing the remaining congruence by 5 gives

$$b + 2c \equiv b' + 2c' \pmod{16}.$$

The four possible translates of $2D$ modulo 16 are

b	$b + 2D \pmod{16}$
0	$\{0, 2, 8, 10\}$
1	$\{1, 3, 9, 11\}$
4	$\{4, 6, 12, 14\}$
13	$\{5, 7, 13, 15\}$

They partition all sixteen residues, and the four residues of $2D$ are distinct. Hence $b = b'$ and $c = c'$. This proves injectivity without relying on an exhaustive computer search. $\square$

[write] **7 October 2026; sources 60–69.** *Checked at the write:* the 64 sums $a+5b+10c$, $a, b, c \in \{0, 1, 4, 13\}$, are distinct modulo 80; $\log_4 80 = 3.1609640474$ and $5 + 9/\log_4 80 = 7.8472326369$.

The seed packs $4^3 = 64$ distinct ordered weighted sums into eighty residues. The accompanying certificate lists all sixty-four images, but the preceding proof is the certificate that matters mathematically.

III.101.3 Why carries do not invalidate digit lifting

Theorem III.101.5 (Cyclic digit lifting). *Let $B \geq 2$ and let $D \subseteq \{0, \dots, B-1\}$ be w -separating modulo B . For $t \geq 0$ define*

$$S_t(D, B) = \left\{ \sum_{j=0}^{t-1} d_j B^j : d_j \in D \right\}, \quad S_0 = \{0\}. \quad (\text{III.101.2})$$

Then $|S_t| = |D|^t$, and S_t is w -separating modulo B^t and over the integers. Every subset of S_t has the same separation property.

Proof. Cardinality follows from uniqueness of base- B expansions. The case $t = 0$ is immediate. For the induction step write each coordinate as $d_j + Bu_j$ and each coordinate on the other side as $d'_j + Bu'_j$. An equality of weighted sums modulo B^t , reduced modulo B , implies $d_j = d'_j$ for all j . Their weighted contributions are consequently *equal as integers*, not merely congruent. Subtract them and divide by B . The remaining equality modulo B^{t-1} gives $u_j = u'_j$ by induction. An integer equality implies the corresponding modular equality, and restriction to a subset preserves injectivity. $\square$

This argument permits carries in individual weighted sums. It does not assume that all weighted digit sums are smaller than B . Indeed the largest weighted digit sum in Proposition III.101.4 is $16 \cdot 13 = 208 > B$. What prevents ambiguity is that equality of the residues determines the *entire tuple* of lowest digits before the next digit is considered.

Corollary III.101.6. *Let $b_6 = \log_4 80$. For every $r \geq 1$ there is a set S of r nonnegative integers with*

$$\max S < \frac{1040}{79} r^{b_6}$$

that is $(1, 5, 10)$ -separating. Its reduction modulo any integer $m > 211r^{b_6}$, and in particular any $m > 256r^{b_6}$, remains separating.

Proof. Set $t = \lceil \log_4 r \rceil$ and take any r elements of $S_t(D, 80)$. For $r = 1$ the set $\{0\}$ suffices. For $r > 1$ we have $80^t < 80r^{b_6}$ and

$$\max S_t = 13 \frac{80^t - 1}{79} < \frac{1040}{79} r^{b_6}.$$

Two weighted sums with weights $(1, 5, 10)$ lie between zero and $16 \max S$. Their difference has absolute value at most $16 \max S$, which is smaller than

$$\frac{16640}{79} r^{b_6} < 211 r^{b_6} < m.$$

A congruence modulo m is therefore an integer equality. The integer separation from Theorem III.101.5 finishes the proof. $\square$

Corollary III.101.7 (Interval growth). *Let $M_6(X)$ be the largest cardinality of a subset of $\{1, \dots, X\}$ with no nontrivial six-binomial solution. For every integer $X \geq 1$,*

$$\frac{1}{4} X^{1/b_6} \leq M_6(X) \leq (16X)^{1/3}. \quad (\text{III.101.3})$$

In particular $3 \leq \beta_6 \leq b_6$.

Proof. Take $t = \lfloor \log_{80} X \rfloor$. The translate $1 + S_t(D, 80)$ lies in $\{1, \dots, X\}$ and has $4^t \geq X^{1/b_6}/4$ elements. Translation preserves the invariant equation. For the upper bound, if $|S| = r$, the r^3 sums $a + 5b + 10c$ are distinct by Lemma III.101.2; they lie in the integer interval $[16, 16X]$. Thus $r^3 \leq 16X - 15 \leq 16X$. $\square$

III.101.4 A CRT family behind the exponent

The construction is not confined to a single six-variable equation.

Theorem III.101.8 (A family for $(1, a, 2a)$). *Let a be a positive odd integer and let $r = 2^t \leq a$, with $t \geq 1$. There exists an r -element alphabet that is $(1, a, 2a)$ -separating modulo $B = ar^2$. Consequently there are integer separating sets of growth $\gg X^{1/b}$ with*

$$b = \log_r(ar^2) = 2 + \log_r a.$$

Proof. Let $E \subseteq \mathbb{Z}/4^t\mathbb{Z}$ consist of the base-4 strings of length t with digits 0 or 1. The map $(u, v) \mapsto u + 2v$ on E^2 is bijective modulo 4^t : in each base-4 digit the two binary choices encode 0, 1, 2, 3, without carry. Give the r elements of E distinct residues modulo a . Since $\gcd(a, 4^t) = 1$, the Chinese remainder theorem lifts these assignments to r residues D modulo $a4^t$.

In a congruence $x + ay + 2az = x' + ay' + 2az'$ modulo $a4^t$, reduction modulo a determines $x = x'$ in D . Division by a then determines y, z from the bijectivity modulo 4^t . Apply Theorem III.101.5 and the interval argument of Corollary III.101.7. $\square$

For $a = 5, r = 4$ this yields the modulus 80. The particularly short integer representatives in Proposition III.101.4 use another valid inner factorization modulo 16; the general CRT assignment need not return those same four representatives. More generally, an inner separating alphabet modulo Q can be lifted to weights $(1, aw_1, \dots, aw_h)$ modulo aQ whenever $\gcd(a, Q) = 1$ and its size is at most a .

III.102 Why a perfect cyclic six-variable alphabet is impossible

Source 64, delivered Section 3.

The lower exponent 3 in Corollary III.101.7 suggests looking for a seed with $B = r^3$. The next argument rules out *every* such cyclic seed for $(1, 5, 10)$, not just the small seeds examined computationally.

Theorem III.102.1 (A valuation obstruction to perfect weighted codes). *Let $w_1, \dots, w_h$ be positive integers. Suppose $D \subseteq \mathbb{Z}/B\mathbb{Z}$ has cardinality $r > 1$ and the weighted-sum map $D^h \rightarrow \mathbb{Z}/B\mathbb{Z}$ is bijective. Then $B = r^h$, and, for every prime $p \mid r$, the h valuations*

$$v_p(w_1), \dots, v_p(w_h)$$

are pairwise distinct.

Proof. The cardinality identity is immediate. Choose representatives of D in $\{0, \dots, B-1\}$ and write its mask polynomial as $A(X) = \sum_{d \in D} X^d$, so $A(1) = r$. Fix $p \mid r$, set $v = v_p(r)$ and $e = hv = v_p(B)$, and let E be the set of $j \in \{1, \dots, e\}$ for which the cyclotomic polynomial Φ_{p^j} divides A in $\mathbb{Z}[X]$. The distinct monic cyclotomic factors are relatively prime over $\mathbb{Q}$, so $\prod_{j \in E} \Phi_{p^j}$ divides A in $\mathbb{Z}[X]$. Evaluating at 1 and using $\Phi_{p^j}(1) = p$ gives

$$|E| \leq v. \quad (\text{III.102.1})$$

For $1 \leq t \leq e$, let ζ be a primitive p^t -th root of unity. Bijectivity of the weighted sum, and $\zeta^B = 1$, give

$$\prod_{i=1}^h A(\zeta^{w_i}) = \sum_{n=0}^{B-1} \zeta^n = 0. \quad (\text{III.102.2})$$

Thus some factor vanishes. It cannot have $v_p(w_i) \geq t$, since that would be $A(1) = r$. Otherwise ζ^{w_i} has exact order $p^{t-v_p(w_i)}$; irreducibility of the corresponding cyclotomic polynomial implies $t - v_p(w_i) \in E$.

Put $V = \{v_p(w_i) : 1 \leq i \leq h\}$. We have proved the covering $\{1, \dots, e\} \subseteq E + V$. Therefore

$$hv = e \leq |E||V| \leq v|V| \leq vh.$$

Every inequality must be an equality, in particular $|V| = h$, as required. $\square$

Corollary III.102.2. *For $w = (1, 5, 10)$ no non-singleton cyclic alphabet is perfect: if the weighted-sum map is injective on an r -element alphabet modulo B , with $r > 1$, then $B \geq r^3 + 1$.*

Proof. Injectivity gives $B \geq r^3$. At equality it gives bijectivity. For $p = 2$ the weight valuations are $(0, 0, 1)$; for $p = 5$ they are $(0, 1, 1)$; for every other prime they are $(0, 0, 0)$. None are pairwise distinct, contrary to Theorem III.102.1 for any prime divisor of r . $\square$

Remark III.102.3 (The limitation of the obstruction). This does not prove $\beta_6 > 3$. A sequence of increasingly large alphabets could have $\log B / \log r \downarrow 3$, and nondigital constructions are not restricted by this theorem. It does show that a fixed finite menu of cyclic seeds cannot attain exponent 3: mixed-radix concatenation has exponent a weighted average of the finitely many numbers $\log B_j / \log r_j > 3$. For comparison, the perfect four-variable seed $(B, D) = (9, \{0, 1, 2\})$ is consistent with the criterion, since $v_3(1), v_3(3)$ are distinct.

The proof gives slightly more information. In a perfect code, for each $p \mid r$ the sets $E + v_p(w_i)$ must partition the interval $\{1, \dots, hv_p(r)\}$, with no overlaps and no values outside that interval. This is a finite necessary condition that can be checked before a seed search. It is not asserted to be sufficient.

III.103 An exact nine-copy circle transfer

Source 64, delivered Section 4.

Fix an even integer $k \geq 4$. A coloring of an interval is *symmetric-progression-free* if every integer progression $b_0, \dots, b_{k-1}$ with $\phi(b_i) = \phi(b_{k-1-i})$ for all i is constant. The order of the colors is irrelevant; only mirror equality is tested.

III.103.1 Three equal arcs prevent a nonconstant palindromic itinerary

Lemma III.103.1 (Equal-arc rigidity). *Partition $\mathbb{R}/(qL\mathbb{Z})$ into its q consecutive half-open arcs of length L , where $q \geq 3$ and $L > 0$. If an even-length arithmetic progression has the same arc label at every pair of reflected positions, then all its points lie in one arc, and their representatives in that arc form a real arithmetic progression.*

Proof. The two central points lie in one arc. Their difference therefore has a representative $\delta \in (-L, L)$. After reversing the progression if necessary, assume $0 \leq \delta < L$. Start from the central pair and extend outwards, one reflected pair at a time. As long as the two previous points lie in the central arc, the next left point can lie only in that arc or its predecessor, and the next right point can lie only in that arc or its successor. The predecessor and successor have different labels because $q \geq 3$. Equality of the reflected labels forces both new points to stay in the central arc. Induction proves the first assertion. All actual consecutive differences in that arc lie in $(-L, L)$ and are congruent to δ modulo qL ; hence they equal δ , giving the second assertion. $\square$

Two arcs do not suffice for this lemma: on $\mathbb{R}/10\mathbb{Z}$ the progression $4, 6, 8, 0$ has the palindromic arc labels $0, 1, 1, 0$ for arcs of length 5. This is an obstruction to that auxiliary lemma, not a proof that nine copies are optimal for every possible transfer construction.

III.103.2 The exact statement

Theorem III.103.2 (Nine-copy transfer). *Let $\phi : \{0, \dots, N-1\} \rightarrow \{1, \dots, r\}$ be symmetric-progression-free. There is a half-open interval coloring $\Phi : \mathbb{T} \rightarrow \{1, \dots, 9r\}$ such that*

$$\mathbb{P}_{x,d \in \mathbb{T}}(\Phi(x+id) = \Phi(x+(k-1-i)d) \text{ for all } i) = \frac{1}{9(k-1)N}. \quad (\text{III.103.1})$$

More precisely, the symmetry event holds if and only if all k points lie in the same member of a partition into $9N$ equal half-open intervals.

Proof. Use indices $a, c \in \{0, 1, 2\}$, $b \in \{0, \dots, N-1\}$, and intervals

$$I_{a,b,c} = \left[\frac{3aN + 3b + c}{9N}, \frac{3aN + 3b + c + 1}{9N} \right). \quad (\text{III.103.2})$$

For $x \in I_{a,b,c}$ set $\Phi(x) = (a, c, \phi(b))$, using any fixed encoding of these triples by $9r$ colors.

Suppose $x_i = x + id$ has symmetric colors. Write

$$3Nx_i = j_i + \theta_i, \quad 0 \leq j_i < 3N, \quad 0 \leq \theta_i < 1, \quad j_i = a_iN + b_i.$$

Then $c_i = \lfloor 3\theta_i \rfloor$, and both a_i and c_i have mirror symmetry. The θ_i form a progression modulo 1. Applying Lemma III.103.1 to three arcs of length $1/3$ shows that all c_i agree and that the chosen θ_i form a real progression, say $\theta_i = \theta_0 + i\delta$. It follows that j_i form an arithmetic progression modulo $3N$: the common difference is congruent to $3N\delta - \delta$, and is an integer because all j_i are integers.

Apply Lemma III.103.1 a second time, now to three arcs of length N on $\mathbb{R}/(3N\mathbb{Z})$. All a_i agree, and j_i are an ordinary integer progression in that block. Thus $b_i = j_i - a_iN$ form an integer progression inside $\{0, \dots, N-1\}$. Their ϕ -colors are symmetric, so the hypothesis forces all b_i to agree. Hence all three cell indices agree. The converse, that points in one cell have symmetric colors, is immediate.

It remains to compute the measure. A cell has length $\ell = 1/(9N) < 1/2$. If all the points of a progression lie in that cell, their real representatives form a real progression, since consecutive differences in $(-\ell, \ell)$ are uniquely determined modulo 1. Its step satisfies $|d| \leq \ell/(k-1)$, and for a fixed such step the admissible starting points have measure $\ell - (k-1)|d|$. Thus the probability for a particular cell is

$$\int_{-\ell/(k-1)}^{\ell/(k-1)} (\ell - (k-1)|d|) \, dd = \frac{\ell^2}{k-1}.$$

Sum over the $9N$ disjoint cells to obtain (III.103.1). All endpoint issues concern null sets and do not change the integral. $\square$

This is a specialized strengthening of the circle-lifting step used in [37, Lemma 4.3]. The specialization is important: the event controlled here is *full mirror symmetry*, not the broader collection of binomial color patterns. The construction uses nine copies for every even k , and the symmetric event is identified exactly rather than estimated by a collision union bound.

III.104 Exact binomial transfer to the two-torus

Source 64, delivered Section 5.

III.104.1 The reduced progression functional

For even $k \geq 4$ let

$$c_i = (-1)^i \binom{k-1}{i} \quad (0 \leq i < k), \quad W_k = \sum_{c_i > 0} c_i = 2^{k-2}.$$

In particular $c_0 = 1$, $c_{k-1} = -1$, $c_{k-1-i} = -c_i$, and $\sum_i c_i = 0$. For a bounded Riemann integrable $F : \mathbb{T}^2 \rightarrow \mathbb{R}$ define

$$\tilde{\Lambda}_k(F) = \int_{\mathbb{T}^2} \int_{\mathbb{T}^{k-1}} \prod_{i=0}^{k-1} F(x + id, y_i) \, dy_0 \cdots dy_{k-2} \, dx \, dd, \quad y_{k-1} = \sum_{i=0}^{k-2} c_i y_i \pmod{1}. \quad (\text{III.104.1})$$

The last coordinate is solved with coefficient -1 , so this is normalized Haar measure on the binomial constraint, with no extra Jacobian or covering multiplicity. This is the reduced functional of [37, Proposition 2.2].

Call $S \subseteq \mathbb{Z}/m\mathbb{Z}$ *mirror-separating* for k if

$$\sum_{i=0}^{k-1} c_i s_i = 0 \text{ in } \mathbb{Z}/m\mathbb{Z}, \quad s_i \in S \quad \implies \quad s_i = s_{k-1-i} \text{ for all } i.$$

This is stronger than merely forbidding nontrivial solutions when additional trivial partition types exist. The distinction is built into the hypothesis.

Define the slice constant

$$V_k = \int_{[0,1]^{k-1}} \mathbf{1}_{[0,1]} \left(\sum_{i=0}^{k-2} c_i t_i \right) dt_0 \cdots dt_{k-2}. \quad (\text{III.104.2})$$

It is positive and at most one. Positivity follows, for example, by taking all t_i sufficiently close to $1/2$: their indicated linear combination is then close to $1/2$, because $\sum_{i < k-1} c_i = 1$.

Theorem III.104.1 (Exact mirror-label transfer). *Let ϕ be an r -coloring of $\{0, \dots, N-1\}$ with no nonconstant symmetrically colored k -progression. Suppose $S \subseteq \mathbb{Z}/m\mathbb{Z}$ contains $9r$ distinct labels and is mirror-separating for k . For every*

$$0 < \alpha \leq \frac{1}{W_k m} \quad (\text{III.104.3})$$

there is an indicator $F : \mathbb{T}^2 \rightarrow \{0, 1\}$, a finite union of half-open rectangles, with constant first marginal α and

$$\boxed{\tilde{\Lambda}_k(F) = \frac{V_k}{9(k-1)N} \alpha^{k-1}.} \quad (\text{III.104.4})$$

Proof. Use the coloring Φ of Theorem III.103.2, and assign its colors distinct representatives $s_1, \dots, s_{9r} \in \{0, \dots, m-1\}$ from S . Set

$$F(x, y) = \mathbf{1}_{[s_{\Phi(x)}/m, s_{\Phi(x)}/m + \alpha)}(y). \quad (\text{III.104.5})$$

The intervals do not cross 1, because $\alpha \leq 1/m$. Every vertical fiber has length α , giving the marginal assertion.

Suppose a configuration contributes to (III.104.1). Write $y_i = s_{\Phi(x+id)}/m + u_i$, with $0 \leq u_i < \alpha$. The binomial relation gives

$$\frac{1}{m} \sum_i c_i s_{\Phi(x+id)} + \sum_i c_i u_i \in \mathbb{Z}, \quad \left| \sum_i c_i u_i \right| < W_k \alpha \leq \frac{1}{m}.$$

The first summand is an integer multiple of $1/m$. Multiplying the relation by m shows that $m \sum_i c_i u_i$ is an integer of absolute value smaller than one; it is therefore zero. Consequently $\sum_i c_i s_{\Phi(x+id)} \equiv 0 \pmod{m}$. Mirror separation and distinctness of the labels force $\Phi(x+id) = \Phi(x+(k-1-i)d)$ for all i . By Theorem III.103.2, all $x+id$ lie in one cell of (III.103.2).

Conversely, fix such a horizontal progression. All its vertical intervals then have the same left endpoint. Since $\sum_i c_i = 0$, translation by that endpoint cancels out of the constraint. Because

$W_k \alpha \leq 1/m \leq 1$, the constraint on the $u_i \in [0, \alpha)$ is an equality over $\mathbb{R}$, not a nonzero integer congruence. Solve for u_{k-1} and rescale the other coordinates by α . The conditional integral is exactly $V_k \alpha^{k-1}$. The horizontal event has probability $1/(9(k-1)N)$ by Theorem III.103.2, proving (III.104.4). $\square$

Remark III.104.2 (No density rounding is needed). The width in (III.104.5) is the desired real density α , not a prescribed reciprocal of an integer. Thus the construction is already an indicator with exactly that marginal. The only issue in reaching arbitrary small α is to choose a coloring and labels satisfying (III.104.3). This observation will remove the endpoint loss.

III.104.2 The slice constants at lengths four and six

Proposition III.104.3. *The constants in (III.104.2) satisfy*

$$V_4 = \frac{8}{27}, \quad V_6 = \frac{3113}{37500}. \quad (\text{III.104.6})$$

More generally, let $w_1, \dots, w_{k/2}$ be the positive half-binomial weights, let w' be the list containing each of them twice, and set $W = \sum_j w_j$. Then

$$V_k = \frac{1}{(k-1)! \prod_j w_j^2} \sum_{J \subseteq \{1, \dots, k\}} (-1)^{|J|} \left(W - \sum_{j \in J} w'_j \right)_+^{k-1}. \quad (\text{III.104.7})$$

Proof. Regard V_k as the density at zero of $\sum_i c_i U_i$, with independent uniform random variables $U_i \in [0, 1]$. This follows directly by conditioning on all but the last variable, whose coefficient has absolute value one. Replace U_i by $1 - U_i$ at the negative coefficients. The new random variable is $\sum_{j=1}^k w'_j U_j - W$.

For positive weights $v_1, \dots, v_k$, inclusion–exclusion applied to the simplex $\{z_j \geq 0 : \sum_j z_j \leq x\}$, with upper bounds $z_j \leq v_j$, gives

$$\mathbb{P} \left(\sum_j v_j U_j \leq x \right) = \frac{1}{k! \prod_j v_j} \sum_J (-1)^{|J|} \left(x - \sum_{j \in J} v_j \right)_+^k.$$

Differentiate and substitute $x = W$ to obtain (III.104.7). At length four use $(w_1, w_2) = (1, 3)$. At length six use $(1, 5, 10)$. In the latter case the numerator sum is 24904 and the denominator is $5!(1 \cdot 5 \cdot 10)^2 = 300000$, giving $3113/37500$. The included exact-arithmetic verifier independently evaluates both finite sums. $\square$

For example, at length six the construction has the exact value

$$\tilde{\Lambda}_6(F) = \frac{3113}{1687500} \frac{\alpha^5}{N}. \quad (\text{III.104.8})$$

This identity isolates the two costs cleanly: N^{-1} comes from the horizontal coloring, and α^5 from the codimension-one vertical slice.

III.105 Endpoint exponents from a prime-indexed family

Source 64, delivered Section 6.

III.105.1 A scale-uniform transfer theorem

Theorem III.105.1 (Endpoint extraction). *Fix even $k \geq 4$, an integer $d \geq 1$, a positive integer c , and real constants $C \geq 1$, $b > 0$. Suppose:*

- (a) *for every prime $p > k$, there is a coloring of $\{0, \dots, p^d - 1\}$ with at most cp colors and no nonconstant symmetrically colored k -progression;*
- (b) *for every integer $R \geq 1$, there is an R -element mirror-separating label set for k modulo some $m_R \leq CR^b$.*

Put

$$A_* = W_k C (9c)^b, \quad \alpha_* = [A_* (2(k+1))^b]^{-1}, \quad K_* = \frac{2^d V_k}{9(k-1)} A_*^{d/b}. \quad (\text{III.105.1})$$

Then, for $0 < \alpha \leq \alpha_$, an indicator function F with constant first marginal α satisfies*

$$\tilde{\Lambda}_k(F) \leq K_* \alpha^{k-1+d/b}. \quad (\text{III.105.2})$$

Consequently $\rho_k(\alpha) \leq K_ \alpha^{k-1+d/b}$.*

Proof. Let $T = (A_* \alpha)^{-1/b}$, so $T \geq 2(k+1)$. The Bertrand–Chebyshev prime interval theorem gives a prime p with $T/2 \leq p \leq T$; this standard real endpoint version follows by applying the integer theorem to $\lceil T/2 \rceil - 1$. In particular $p > k$. Pad the coloring palette, when necessary, to exactly cp colors. Take $R = 9cp$ labels and their modulus m_R . Then

$$W_k m_R \alpha \leq W_k C (9cp)^b \alpha = A_* p^b \alpha \leq 1.$$

Apply Theorem III.104.1 with $N = p^d$ and the *actual* width α . Since $p \geq T/2$,

$$\tilde{\Lambda}_k(F) = \frac{V_k}{9(k-1)p^d} \alpha^{k-1} \leq \frac{2^d V_k}{9(k-1)} A_*^{d/b} \alpha^{k-1+d/b}.$$

The passage to ρ_k is the analytic realization in Section III.106 (or the corresponding published realization theorem in [37]). All constants were chosen before α ; there is no limiting choice of a fixed prime hidden in the implied constant. $\square$

A useful more flexible version replaces p^d by a length at least $c_0 p^d$ and cp by any palette bounded by a fixed multiple of p . Only the constants change. The essential features are scale-uniform polynomial label size and a coloring at a prime within a fixed multiplicative factor of the target scale.

III.105.2 The external coloring input

The input from Shi–Dong is the following theorem [33, Theorem 1.1]: for even $k \geq 4$ and prime $p > k$ there is a $(k-1)^{k^2/4-1} p$ -coloring of $\mathbb{Z}/p^{k^2/4}\mathbb{Z}$ without nonconstant symmetrically colored k -progressions. Restriction gives the required interval coloring. Their construction combines carry-control colors on base- p digits with a layered field-norm polynomial. We use the theorem with

$$d = k^2/4, \quad c = (k-1)^{d-1}. \quad (\text{III.105.3})$$

This input is a credited 2026 preprint theorem, not a new construction asserted by the present manuscript. The exact degree-restricted barrier proved in Section III.107 clarifies one of its limitations.

III.105.3 Length four

The alphabet $D_4 = \{0, 1, 2\}$ modulo 9 is $(1, 3)$ -separating: the sum $a + 3b$ is the ordinary two-digit base-3 representation of an integer in $\{0, \dots, 8\}$. By Theorem III.101.5, for any R we may select R elements of $S_t(D_4, 9)$, where $t = \lceil \log_3 R \rceil$, and use the modulus 9^t . It satisfies $9^t \leq 9R^2$. Thus Theorem III.105.1 applies with

$$k = 4, \quad d = 4, \quad c = 27, \quad b = 2, \quad C = 9, \quad (\text{III.105.4})$$

and gives $\rho_4(\alpha) \ll \alpha^5$. The seed and its exponent are the established four-variable construction; the endpoint extraction is the improvement.

III.105.4 Length six

Select R elements of $S_t(D, 80)$ with $t = \lceil \log_4 R \rceil$. The modulus 80^t itself works by Theorem III.101.5 and satisfies $80^t \leq 80R^{b_6}$. This is more efficient for the present application than the all-moduli bound in Corollary III.101.6: here we are free to choose the modulus. Use

$$k = 6, \quad d = 9, \quad c = 5^8, \quad b = b_6, \quad C = 80. \quad (\text{III.105.5})$$

Then Theorem III.105.1 gives

$$\rho_6(\alpha) \ll \alpha^{5+9/b_6} = \alpha^{7.847232636916081\dots}.$$

This proves part (ii) of Theorem F.69.1, once the deterministic realization below is supplied.

Case	d	b	Count exponent
4	4	2	5
6 with generic labels	9	5	34/5 before losses
6 in this manuscript	9	$\log_4 80$	7.8472326369...
6 with conjecturally optimal-order labels	9	3	8

The last line is conditional on the stronger uniform estimate $m_R \leq CR^3$, not merely on the definition-level statement $\beta_6 = 3$. An $R^{3+o(1)}$ label bound by itself only yields the corresponding exponent $8 - o(1)$ through this argument.

What the bounds do not say. They do not say that the implied constants are small. The explicit constants in (III.105.1) are large and were not optimized. They do not provide a useful finite threshold for the final prime modulus as a function of density and uniformity error. They do show a genuine fixed-power improvement as $\alpha \downarrow 0$, not a bound with a constant secretly varying with α .

III.106 Deterministic realization in cyclic groups

Source 64, delivered Section 7.

For completeness, we prove the prime-modulus version of the realization needed here. The general torus realization is an established result [37, Proposition 2.2]; the proof below makes the dependence and the indicator nature of our examples explicit. No probabilistic rounding is used.

Theorem III.106.1 (Prime-modulus realization). *Fix $k \geq 4$ and a Riemann integrable indicator $F : \mathbb{T}^2 \rightarrow \{0, 1\}$ with constant first marginal α . For primes P set*

$$A_P = \left\{ n \in \mathbb{Z}/P\mathbb{Z} : F\left(\frac{n}{P} \bmod 1, \frac{n^{k-2}}{P} \bmod 1\right) = 1 \right\}. \quad (\text{III.106.1})$$

Then, as $P \rightarrow \infty$ through primes,

$$\frac{|A_P|}{P} \rightarrow \alpha, \quad \|\mathbf{1}_{A_P} - \alpha\|_{U^{k-2}} \rightarrow 0, \quad \Lambda_k(\mathbf{1}_{A_P}) \rightarrow \tilde{\Lambda}_k(F). \quad (\text{III.106.2})$$

The prime P in this theorem is distinct from the prime p selected in Theorem III.105.1. First α determines p , the finite coloring, the labels and F . Only then does $P \rightarrow \infty$, with all of those objects held fixed.

III.106.1 An elementary complete-sum estimate

Write $e_P(u) = \exp(2\pi i u/P)$ for $u \in \mathbb{F}_P$.

Lemma III.106.2. *Let $H \in \mathbb{F}_P[X_1, \dots, X_m]$ have degree $s \geq 2$, with $P > s$ and a nonzero homogeneous part of degree s . Then*

$$\left| \mathbb{E}_{z \in \mathbb{F}_P^m} e_P(H(z)) \right| \leq \left(\frac{s-1}{P} \right)^{1/2^{s-1}}.$$

A nonconstant linear polynomial has complete average zero.

Proof. There is a vector v on which the top homogeneous part is nonzero, since a nonzero polynomial of total degree less than P cannot vanish at every point of $\mathbb{F}_P^m$. The latter assertion follows by induction on the number of variables from the univariate root bound. Extend v to a basis and average first in its coordinate t . For each choice of the other coordinates the resulting polynomial in t has degree s with nonzero leading coefficient.

For such a polynomial Q , repeated Cauchy–Schwarz, equivalently $|\mathbb{E}f| \leq \|f\|_{U^{s-1}}$, bounds $|\mathbb{E}_t e_P(Q(t))|^{2^{s-1}}$ by the average of the $(s-1)$ -fold multiplicative derivative. Its exponent is linear in t , with coefficient $s!a_s h_1 \cdots h_{s-1}$. When every h_j is nonzero its complete t -average is zero; when some h_j is zero the multiplicative derivative is identically one. The average is therefore $1 - (1 - 1/P)^{s-1} \leq (s-1)/P$. Take roots and then average over the remaining coordinates. The linear assertion is character orthogonality. $\square$

III.106.2 The progression distribution

Put $s = k - 2$. The integer polynomials $(X + iY)^s$ for $0 \leq i < k$ span the space of homogeneous degree- s polynomials in two variables. Their only linear relation is a scalar multiple of

$$\sum_{i=0}^{k-1} (-1)^i \binom{k-1}{i} (X + iY)^s = 0. \quad (\text{III.106.3})$$

Indeed the coefficients are a Vandermonde matrix of rank $s + 1 = k - 1$, and (III.106.3) is the $(k-1)$ -st finite difference of a polynomial of smaller degree.

Consider the empirical distribution on $\mathbb{T}^{2k}$ of

$$\left(\frac{n+id}{P} \bmod 1, \frac{(n+id)^s}{P} \bmod 1 \right)_{i=0}^{k-1}, \quad n, d \in \mathbb{F}_P.$$

Test it against an arbitrary character with integer frequencies r_i, t_i . Its average is the complete exponential sum for

$$\sum_i r_i(n+id) + \sum_i t_i(n+id)^s.$$

If the degree- s polynomial is nonzero over $\mathbb{Q}$, it remains nonzero modulo all sufficiently large primes and Lemma III.106.2 makes the average tend to zero. If it vanishes identically, the remaining average is zero unless $\sum_i r_i = \sum_i i r_i = 0$, in which case it is one for all sufficiently large primes. These are exactly the annihilating characters of the torus on which the first coordinates are an arithmetic progression and the second coordinates satisfy the single binomial equation. The character criterion for weak convergence therefore gives Haar measure on that torus.

Each coordinate projection onto a pair (x_i, y_i) has uniform measure on $\mathbb{T}^2$. A Riemann integrable F has a discontinuity set of two-dimensional measure zero; the product of the k copies of F is consequently continuous almost everywhere for the limiting Haar measure. Approximation by continuous functions, or the standard almost-everywhere-continuity form of weak convergence, now gives the last limit in (III.106.2). The same argument with a single pair gives the density limit.

III.106.3 The cube distribution and centered uniformity

We need one further elementary independence fact:

$$\{(X + \omega \cdot H)^s : \omega \in \{0, 1\}^s\} \quad \text{is linearly independent over } \mathbb{Q}. \quad (\text{III.106.4})$$

To prove this, suppose coefficients a_ω give a zero combination. For $I \subseteq \{1, \dots, s\}$, the coefficient of $X^{s-|I|} \prod_{j \in I} H_j$ is

$$\frac{s!}{(s-|I|)!} \sum_{\omega: \omega_j=1 \ (j \in I)} a_\omega.$$

Starting at $I = \{1, \dots, s\}$ and descending in cardinality forces every a_ω to vanish.

Apply the character argument just used to the 2^s cube pairs

$$\left(\frac{n + \omega \cdot h}{P}, \frac{(n + \omega \cdot h)^s}{P} \right)_{\omega \in \{0,1\}^s} \pmod{1}.$$

By (III.106.4), their limiting first coordinates form a Haar cube, while *all* second coordinates are independent Haar variables, also independent of the first coordinates. For the cube expression of $F - \alpha$, integration in any second coordinate gives zero, since $\int_{\mathbb{T}} (F(x, y) - \alpha) dy = 0$ for almost every x . Thus the limiting cube moment is zero, proving the middle limit of (III.106.2) and completing Theorem III.106.1.

The deduction of an upper bound on ρ_k . If the right side of (F.69.3) had a lower-bound value strictly larger than $\tilde{\Lambda}_k(F)$, choose ε smaller than half the gap. The sequence in Theorem III.106.1 would violate that lower bound for arbitrarily large prime P . Therefore $\rho_k(\alpha) \leq \tilde{\Lambda}_k(F)$, with the quantifiers in the definition fully respected.

III.106.4 Exact rational density, rather than merely limiting density

Corollary III.106.3. *Suppose $\alpha = a/q \in (0, 1)$ is rational, and F is as in Theorem III.106.1. There are deterministic sets $B_P \subseteq \mathbb{Z}/(qP)\mathbb{Z}$ with density exactly α , centered U^{k-2} norm tending to zero, and $\Lambda_k(\mathbf{1}_{B_P}) \rightarrow \tilde{\Lambda}_k(F)$ as $P \rightarrow \infty$ through primes.*

Proof. Pull A_P back under the quotient map $\mathbb{Z}/(qP)\mathbb{Z} \rightarrow \mathbb{Z}/P\mathbb{Z}$. All normalized cube and progression averages are unchanged by this pullback, since the quotient map is surjective and each fiber has equal cardinality. Its density is $|A_P|/P \rightarrow \alpha$. Add or remove the smallest available residues, in their standard order, until the cardinality is aP . This changes a proportion $\delta_P = ||A_P|/P - \alpha| \rightarrow 0$ of the group. For the difference g of the two indicators, boundedness and the cube formula give $\|g\|_{U^s}^{2^s} \leq \mathbb{E}|g| = \delta_P$. The Gowers norm triangle inequality preserves the centered uniformity limit. A telescoping expansion of the progression product changes Λ_k by at most $k\delta_P$. $\square$

This corollary matters for formulations that fix the density before choosing the ambient modulus. It avoids inferring an exact-density statement merely from a sequence of slightly varying densities.

III.107 An exact dimension barrier for bounded-degree polynomial colors

Source 64, delivered Section 8.

The coloring dimension $k^2/4$ enters Theorem III.105.1 in the numerator of the improvement d/b . A tempting way to improve the final exponent would be to increase that dimension while retaining a single degree- $(k-1)$ polynomial color. Within that model this is impossible.

Lemma III.107.1 (Chevalley–Warning in the needed form). *Let $P_1, \dots, P_t \in \mathbb{F}_p[X_1, \dots, X_n]$ vanish at the origin. If $\sum_j \deg P_j < n$, then they have a common nonzero zero in $\mathbb{F}_p^n$. Zero polynomials may be omitted from the degree sum.*

Proof. Modulo p , the number of common zeros is

$$\sum_{x \in \mathbb{F}_p^n} \prod_{j=1}^t (1 - P_j(x)^{p-1}).$$

The summand has degree less than $n(p-1)$. Every monomial of this degree has some variable whose exponent is less than $p-1$, and its sum over that variable is zero in $\mathbb{F}_p$ (also for exponent zero, when the sum is p). Thus the number of common zeros is divisible by p . It is positive since the origin is a zero, and therefore at least p . $\square$

Theorem III.107.2 (Exact degree-restricted dimension barrier). *Let $k = 2s \geq 4$, let $p > k$ be prime, and let $Q = (Q_1, \dots, Q_t) : \mathbb{F}_p^n \rightarrow \mathbb{F}_p^t$ have polynomial coordinates of total degree at most $k-1$. Suppose the coloring by Q has no nonconstant symmetrically colored k -progression. Then*

$$n \leq ts^2 = \frac{tk^2}{4}. \quad (\text{III.107.1})$$

For a scalar color ($t = 1$), the bound is sharp.

Proof. Fix any center $a \in \mathbb{F}_p^n$. Expand

$$Q_j(a + Th) = \sum_{r=0}^{2s-1} B_{j,r}(h)T^r.$$

The coefficient $B_{j,r}(h)$ is homogeneous of degree r in h , possibly zero. If $n > ts^2$, the system

$$B_{j,1}(h) = B_{j,3}(h) = \cdots = B_{j,2s-1}(h) = 0 \quad (1 \leq j \leq t)$$

has total degree at most $t(1+3+\cdots+(2s-1)) = ts^2 < n$ and vanishes at $h = 0$. By Lemma III.107.1 it has a nonzero solution. For this h , every $Q_j(a + Th)$ is even in T . The k points

$$a + \left(i - \frac{k-1}{2}\right)h, \quad 0 \leq i < k,$$

are a nonconstant progression, and their Q -colors have mirror symmetry. This contradicts the hypothesis and proves (III.107.1).

For sharpness when $t = 1$, use blocks of dimensions $1, 3, \dots, 2s-1$, with anisotropic homogeneous norm forms N_r of degree r , and let $Q = \sum_{r \text{ odd}} N_r$. If a line has symmetric Q -values at the k positions, its centered odd part has k distinct roots and degree at most $k-1$, so it vanishes identically. Descending through the odd degrees, the highest remaining coefficient is $N_r(h_r)$ and forces $h_r = 0$. This is the layered norm stage of the Shi–Dong construction, and it has exactly s^2 coordinates. $\square$

[write] 7 October 2026; sources 60–69. *This is source 10's exact threshold* (Theorem III.30.3, Section III.30) at the degree bounds $D_\ell = k-1$: there $s_\ell = \lceil (k-1)/2 \rceil = k/2$ and $B = \sum_\ell s_\ell^2 = tk^2/4$, and source 10 proves both the bound and its sharpness for every t , by the same Chevalley–Warning argument; source 14 reproves it (Theorem III.36.1). This source cites neither; both were in the report at its search snapshots (checked at the write). It is kept as printed, a second proof of a special case.

The barrier does not cover arbitrary colorings, higher-degree polynomials, nonpolynomial coordinate maps, or a construction whose number of colors grows in another way. It sharpens the degree- $(k-1)$ instance of the polynomial-method limitation discussed in [33, Remark 4.2], not every possible approach to symmetric-progression-free colorings.

III.108 Source 67's Section 6: norm layers without a degree cap [write]

Added in the eleventh write, 7 October 2026.

The next section is source 67's Section 6. *Logically it continues source 22's chapter* on polynomial colourings (Section III.60), source 10's dimension threshold and source 38's six-term norm defects (Section III.87); it is printed here because no number printed earlier may move.

What is proved. For q of characteristic $p > 2m$ and the map $c + \sum_j v_j N_{\mathbb{F}_{q^{d_j}}/\mathbb{F}_q}(x_j)$ on the full product of its norm blocks, with $v_j \in \mathbb{F}_q^r$: when $m \in \{2, 3, 4\}$, or $q \geq (2m-1)^2$, the map avoids every nontrivial symmetrically coloured $2m$ -term progression exactly when every d_j is odd and below $2m$ and the coefficient vectors of each fixed degree are independent (Theorem III.109.1), so the maximum source dimension is rm^2 with no degree cap assumed in advance. The proof counts progressions by character sums (Theorem III.109.3), uses exact positivity at lengths four and six (Theorem III.109.5)

and eleven explicit degree-nine irreducibility certificates at length eight (Theorem III.109.6); for general m only degree $2m + 1$ over finitely many small prime fields remains (Proposition III.109.4). The low-degree construction is Shi and Dong's and the degree-capped optimality is source 10's threshold ($B = \sum_{\ell} s_{\ell}^2$, here rm^2), source 22's and source 38's, as source 67 credits; the new part excludes every higher norm degree. The unrestricted scalar quintic for four-term symmetry stays open (item III.19).

Relation to the formal statements. The colouring problem has no Lean statement; nothing in this chapter is formalized.

III.109 A degree-free obstruction for independent norm layers

Source 67, delivered Section 6.

The field-norm construction of Shi and Dong uses disjoint blocks of odd dimensions $1, 3, \dots, 2m - 1$ to avoid symmetrically coloured $2m$ -term arithmetic progressions [33]. The preceding ProveIt reports determine its optimal dimension among polynomials whose degrees are at most $2m$, and isolate a genuine unresolved question for unrestricted scalar quintics and four-term progressions [272]. We establish a different, complementary barrier: within the full class of sums of independent norm blocks, arbitrarily high degrees do not improve the construction. The result is unconditional for progression lengths 4, 6, and 8, and for all lengths once $q \geq (2m - 1)^2$. The eight-term case uses eleven small exact irreducibility certificates. The general character-sum argument also reduces every remaining degree to the single critical degree $2m + 1$ over finitely many prime fields.

III.109.1 Definitions and the main classification

Let q be a prime power of characteristic $p > 2m$, where $m \geq 2$. Put $\alpha_i = 2i - 1$ for $1 \leq i \leq m$. These elements and their negatives are pairwise distinct: $\alpha_i - \alpha_j = 2(i - j)$ and $\alpha_i + \alpha_j = 2(i + j - 1)$. Here 2 is invertible, $0 < |i - j| < m < p$ for $i \neq j$, and $0 < i + j - 1 \leq 2m - 1 < p$. A map $C : V \rightarrow \mathbb{F}_q^r$ avoids symmetric $2m$ -term progressions if

$$C(z + \alpha_i h) = C(z - \alpha_i h) \quad (1 \leq i \leq m) \implies h = 0.$$

The points in increasing integer-node order are an arithmetic progression with common difference $2h$, so this is the usual notion. We count ordered progressions, or equivalently their pairs (z, h) with $h \neq 0$; reversal is counted separately.

For positive integers d_j , let $K_j = \mathbb{F}_{q^{d_j}}$ and $N_j = N_{K_j/\mathbb{F}_q}$. An *independent norm-block map* is

$$C(x_1, \dots, x_s) = c + \sum_{j=1}^s v_j N_j(x_j), \quad x_j \in K_j, \quad v_j \in \mathbb{F}_q^r. \quad (\text{III.109.1})$$

The direct-product hypothesis on the source is part of this class: the blocks have independent coordinates. Zero coefficient vectors are permitted in the definition and will correctly be excluded by the theorem. Arbitrary invertible affine changes of source coordinates and bijective relabellings of the colours preserve the conclusion.

Theorem III.109.1 (Complete norm-block classification). *Assume either $m \in \{2, 3, 4\}$ or $q \geq (2m - 1)^2$. The map (III.109.1) avoids every nontrivial symmetric $2m$ -term progression if and only if both conditions hold:*

1. Every d_j is odd and satisfies $d_j < 2m$.
2. For each fixed degree d , the vectors $\{v_j : d_j = d\}$ are linearly independent over $\mathbb{F}_q$.

Consequently its source dimension is at most

$$\sum_j d_j \leq r \sum_{i=1}^m (2i-1) = rm^2. \quad (\text{III.109.2})$$

This bound is attained, and requires no a priori degree cap. In particular the exact maximum dimensions are $4r$, $9r$, and $16r$ for lengths 4, 6, and 8, respectively, over every field of characteristic larger than the progression length.

[write] 7 October 2026; sources 60–69. With degree caps this is source 10's threshold: odd degrees $d_j < 2m$ have $s = \lceil d_j/2 \rceil$, and independent blocks reach $\sum s_\ell^2 = rm^2$ (Theorem III.30.3); the content of this theorem is that no higher norm degree helps.

The sufficiency and the example attaining (III.109.2) are the familiar descending-degree argument behind the layered construction. The new obstruction is the exclusion of every higher odd norm degree. It is proved quantitatively below. The theorem concerns (III.109.1); it does not determine the maximum dimension for general scalar quintic polynomials, whose interacting terms need not have this form.

III.109.2 An exact reduction to one norm-ratio problem

Fix $d \geq 1$, write $K = \mathbb{F}_{q^d}$, $Q = q^d$, and $N = N_{K/\mathbb{F}_q}$. Define

$$M_{d,m}(q) = \#\{a \in K : N(a + \alpha_i) = N(a - \alpha_i) \text{ for all } i\}. \quad (\text{III.109.3})$$

None of the $2m$ elements $\pm\alpha_i$ is counted, since at such an element one side of one equality is zero and the other is not. For all other a , the defining conditions are

$$N\left(\frac{a + \alpha_i}{a - \alpha_i}\right) = 1 \quad (1 \leq i \leq m).$$

Proposition III.109.2 (Exact progression and centre counts). *The number of nontrivial ordered symmetric progressions for the single norm $N : K \rightarrow \mathbb{F}_q$ is*

$$A_{2m}^{\text{sym}}(N) = (Q - 1)M_{d,m}(q). \quad (\text{III.109.4})$$

For odd d , the centre zero has no nonzero successful direction, and every nonzero centre has exactly $M_{d,m}(q)$ successful directions. For even d , the corresponding counts are $Q - 1$ at zero and $M_{d,m}(q) - 1$ at every nonzero centre.

Proof. For $h \neq 0$, put $a = z/h$. Multiplicativity gives

$$N(z \pm \alpha_i h) = N(h)N(a \pm \alpha_i),$$

so for each fixed h the admissible centres are precisely h times the set in (III.109.3). This proves (III.109.4). For fixed $z \neq 0$, the map $h \mapsto z/h$ bijects $K^\times$ with $K^\times$. The element $a = 0$ is counted in (III.109.3) exactly when d is even, since $N(-u) = (-1)^d N(u)$. The asserted centre counts follow. $\square$

III.109.3 The character estimate, including its exceptional terms

We use the classical multiplicative Weil bound in its *distinct-root* form [274, 275]: if η has order $e > 1$, $F \in K[X]$ has t distinct roots in an algebraic closure, and F is not a constant times an e th power, then, with $\eta(0) = 0$,

$$\left| \sum_{a \in K} \eta(F(a)) \right| \leq (t-1)\sqrt{Q}. \quad (\text{III.109.5})$$

This external theorem is the analytic input. The application below checks the non-power hypothesis at each active linear factor; polynomial degree with multiplicities is never substituted for t .

For $1 \leq s \leq m$ put

$$b_{m,1}(Q) = 2m-1, \quad b_{m,s}(Q) = (2s-1)\sqrt{Q} + 2(m-s) \quad (s \geq 2).$$

Theorem III.109.3 (Uniform norm-ratio count). *For every $d \geq 1$,*

$$|(q-1)^m M_{d,m}(q) - (Q-2m)| \leq \sum_{s=1}^m \binom{m}{s} (q-2)^s b_{m,s}(Q). \quad (\text{III.109.6})$$

In particular,

$$M_{d,m}(q) \geq \frac{Q-2m - (2m-1)((q-1)^m - 1)\sqrt{Q}}{(q-1)^m}. \quad (\text{III.109.7})$$

Thus, for fixed m, d with $d > 2m$,

$$M_{d,m}(q) = \frac{q^d}{(q-1)^m} + O_m(q^{d/2}), \quad A_{2m}^{\text{sym}}(N) = \frac{q^d(q^d-1)}{(q-1)^m} + O_m(q^{3d/2}). \quad (\text{III.109.8})$$

The error constant is explicit in (III.109.6).

Proof. Let $\widehat{\mathbb{F}_q^\times}$ be the group of multiplicative characters. Orthogonality gives, for $u \in K^\times$,

$$\mathbf{1}_{N(u)=1} = \frac{1}{q-1} \sum_{\chi \in \widehat{\mathbb{F}_q^\times}} \chi(N(u)).$$

Write $E = K \setminus \{\pm\alpha_1, \dots, \pm\alpha_m\}$. Expanding all m indicators proves the exact identity

$$(q-1)^m M_{d,m}(q) = \sum_{\chi_1, \dots, \chi_m} \sum_{a \in E} \prod_{i=1}^m (\chi_i \circ N) \left(\frac{a + \alpha_i}{a - \alpha_i} \right). \quad (\text{III.109.9})$$

The all-principal tuple contributes $Q - 2m$.

Suppose the tuple has s nonprincipal coordinates. Choose a character χ_0 of order $q-1$ on $\mathbb{F}_q^\times$ and write its active characters as $\chi_i = \chi_0^{e_i}$, where $1 \leq e_i \leq q-2$. The character $\eta = \chi_0 \circ N$ also has order $q-1$, because the norm is surjective. On E the summand is $\eta(F(a))$, where

$$F(X) = \prod_{i \text{ active}} (X + \alpha_i)^{e_i} (X - \alpha_i)^{q-1-e_i}.$$

This polynomial has exactly $2s$ distinct roots. Every displayed multiplicity lies strictly between zero and $q-1$, so F is not a $(q-1)$ st power times a constant. Applying (III.109.5) to the sum over all of K , and then removing the $2(m-s)$ inactive forbidden nodes, gives

$$\left| \sum_{a \in E} \eta(F(a)) \right| \leq (2s-1)\sqrt{Q} + 2(m-s).$$

At the active forbidden nodes F vanishes, so no further correction is needed.

For $s = 1$ this can be sharpened without Weil. If the active index is i and its character is $\chi \neq 1$, the substitution $t = (a + \alpha_i)/(a - \alpha_i)$ is a Mobius bijection. Summing the nonprincipal character over $K^\times \setminus \{1\}$ gives -1 . Removing the other $2m - 2$ nodes yields the exact expression

$$-1 - \sum_{j \neq i} \left[\chi \left(\left(\frac{\alpha_j + \alpha_i}{\alpha_j - \alpha_i} \right)^d \right) + \chi \left(\left(\frac{\alpha_j - \alpha_i}{\alpha_j + \alpha_i} \right)^d \right) \right], \quad (\text{III.109.10})$$

whose absolute value is at most $2m - 1$. There are $\binom{m}{s}(q - 2)^s$ character tuples with support size s . The triangle inequality proves (III.109.6). Finally $b_{m,s}(Q) \leq (2m - 1)\sqrt{Q}$ for every s , which proves (III.109.7); dividing by $(q - 1)^m$ and using Proposition III.109.2 proves (III.109.8). $\square$

One can retain the arithmetic content of the singleton terms. For odd d , define

$$\kappa_{d,m}(q) = \# \left\{ (i, j) : i \neq j, \quad \left(\frac{\alpha_j + \alpha_i}{\alpha_j - \alpha_i} \right)^d = 1 \right\}.$$

The ordered pairs (i, j) and (j, i) cannot both be counted, since their ratios are negatives of one another and d is odd. Hence $0 \leq \kappa_{d,m}(q) \leq m(m - 1)/2$. Summing (III.109.10) over its nonprincipal characters gives

$$|(q - 1)^m M_{d,m}(q) - \{Q - mq + 2m(m - 1) - 2(q - 1)\kappa_{d,m}(q)\}| \leq B_m(q)\sqrt{Q} + C_m(q), \quad (\text{III.109.11})$$

where

$$B_m(q) = \sum_{s=2}^m \binom{m}{s} (q - 2)^s (2s - 1), \quad C_m(q) = 2 \sum_{s=2}^m \binom{m}{s} (q - 2)^s (m - s).$$

Indeed, for each active index the sum of its singleton expressions is $2m - q - 2(q - 1)\kappa_i$, by ordinary character orthogonality. This proves (III.109.11) directly.

For four terms it specializes to

$$\left| (q - 1)^2 M_{d,2}(q) - \{q^d - 2q + 4 - 2(q - 1)\kappa\} \right| \leq 3(q - 2)^2 q^{d/2}, \quad \kappa = \mathbf{1}_{2^d \in \{1, -1\}}. \quad (\text{III.109.12})$$

At $d = 5$ and $p > 4$, the correction $\kappa = 1$ occurs precisely in characteristics 11 and 31. This is a correction to the counting estimate, not an exception to the existence result.

III.109.4 Excluding every high degree

Proposition III.109.4 (The only potentially exceptional norm degree). *For every $p > 2m$ the following statements hold.*

1. *If d is even, the norm has a nontrivial symmetric progression centred at zero.*
2. *If d is odd and $d < 2m$, the norm avoids every nontrivial symmetric progression.*
3. *If $d \geq 2m + 3$, then $M_{d,m}(q) > 0$.*
4. *If $d \geq 2m + 1$ and $q \geq (2m - 1)^2$, then $M_{d,m}(q) > 0$.*

Thus outside the explicit field-size range, only the odd degree $d = 2m + 1$ can remain undecided by these statements.

Proof. The first assertion follows from $N(-u) = N(u)$. For the second, the odd polynomial

$$F(T) = N(z + Th) - N(z - Th)$$

has degree at most $d < 2m$. If all reflection equalities hold, it has the $2m$ distinct roots $\pm\alpha_i$, so it is the zero polynomial. But for odd d its leading coefficient is $2N(h)$, which cannot vanish when $h \neq 0$.

For the remaining assertions put $a = (2m-1)((q-1)^m - 1)$ and $t = q^{d/2}$. It suffices, by (III.109.7), to prove $t^2 - at - 2m > 0$. If $q \geq (2m-1)^2$ and $d \geq 2m+1$, then $t \geq (2m-1)q^m$ and consequently

$$t - a \geq (2m-1)\{q^m - (q-1)^m + 1\} \geq 2m-1.$$

Since $t \geq q \geq 2m+1$, we have $t(t-a) > 2m$. If $d \geq 2m+3$ without the extra field-size hypothesis, then $t \geq q^{m+3/2}$, whereas $a < (2m-1)q^m < q^{m+1}$; again $t(t-a) > 2m$. $\square$

Theorem III.109.5 (Four- and six-term barriers over every field). *For $m = 2$ or $m = 3$, $p > 2m$, and every odd $d \geq 2m+1$, one has $M_{d,m}(q) > 0$. More explicitly,*

$$(q-1)^2 M_{d,2}(q) \geq q^d - 3(q-2)^2 q^{d/2} - (6q-8), \quad (\text{III.109.13})$$

$$(q-1)^3 M_{d,3}(q) \geq q^d - (q-2)^2(5q-1)q^{d/2} - 3q(2q-3). \quad (\text{III.109.14})$$

Both right-hand sides are strictly positive throughout their stated ranges. In particular every nonzero centre of a quintic norm has a nontrivial symmetric four-term progression, in every characteristic $p > 4$.

Proof. For $m = 2$, the singleton contribution in (III.109.6) is at most $6(q-2)$, and the support-two contribution at most $3(q-2)^2\sqrt{Q}$. This yields (III.109.13). For $m = 3$, the contributions are at most

$$15(q-2), \quad 3(q-2)^2(3\sqrt{Q}+2), \quad 5(q-2)^3\sqrt{Q},$$

which combine with the 6 excluded nodes to give (III.109.14).

Here are exact positivity certificates; no approximation to $q^{d/2}$ is needed. If $Q > C^2 + 2D$ with $D > 0$, then $Q - D > 0$ and

$$(Q - D)^2 - C^2 Q = Q(Q - C^2 - 2D) + D^2 > 0,$$

so $Q - C\sqrt{Q} - D > 0$. It is enough to check this at the smallest Q in each case, since C, D depend only on q . For four terms put $C = 3(q-2)^2$, $D = 6q-8$, and $q = 5+x$. A direct expansion gives

$$q^5 - C^2 - 2D = x^5 + 16x^4 + 142x^3 + 764x^2 + 2141x + 2352 > 0.$$

For six terms put $C = (q-2)^2(5q-1)$ and $D = 3q(2q-3)$. At $q = 7$ the value of $q^7 - C^2 - 2D$ is $100581 > 0$. Every other prime power of characteristic greater than six is at least eleven. With $q = 11+x$ the same expression is

$$x^7 + 52x^6 + 1101x^5 + 12079x^4 + 72119x^3 + 226218x^2 + 354685x + 354041 > 0.$$

Thus $Q = q^d$ satisfies the required inequality whenever $d \geq 5$ for four terms, or $d \geq 7$ for six terms. The statement about nonzero centres follows from Proposition III.109.2. $\square$

III.109.5 Eleven certificates complete the eight-term case

Theorem III.109.6 (Eight-term norm barrier over every field). *If $p > 8$ and $d \geq 9$ is odd, then $M_{d,4}(q) > 0$.*

Proof. Proposition III.109.4 handles $d \geq 11$, and it handles $d = 9$ when $q \geq 49$. If $p > 8$ and $q < 49$, the field size q must itself be prime, since $p^2 \geq 121$. The possible primes and exact certificates are as follows. Put

$$G(X) = X(X^2 - 1)(X^2 - 9)(X^2 - 25)(X^2 - 49).$$

For each row of the table, $F(X) = G(X) + c$ is irreducible over $\mathbb{F}_p$:

p	11	13	17	19	23	29	31	37	41	43	47
c	1	6	1	6	6	3	3	5	9	4	1

If a is a root in $\mathbb{F}_{p^9}$, then for every $t \in \mathbb{F}_p$

$$N_{\mathbb{F}_{p^9}/\mathbb{F}_p}(a - t) = -F(t).$$

In particular all eight points $a \pm 1, a \pm 3, a \pm 5, a \pm 7$ have norm $-c \neq 0$. These certificates even give monochromatic progressions. The nine conjugates of a and the nine negatives of those conjugates are all successful and are disjoint: the former are roots of $G + c$ and the latter roots of $G - c$, which have no common root since $2c \neq 0$. Thus $M_{9,4}(p) \geq 18$ in every row, giving at least $18(p^9 - 1)$ ordered symmetric progressions.

For completeness, the irreducibility assertion is a finite polynomial certificate with an elementary verification rule. For each displayed pair (p, c) the companion verifier checks

$$X^{p^9} \equiv X \pmod{F}, \quad \gcd(X^{p^3} - X, F) = 1.$$

Every irreducible factor of F therefore has degree dividing nine, and degrees one and three are excluded. Since F has degree nine, it is irreducible. The first congruence also supplies squarefreeness, as $X^{p^9} - X$ has derivative -1 . Independently, the verifier checks $\gcd(X^{p^i} - X, F) = 1$ for $1 \leq i \leq 4$; a reducible polynomial of degree nine would have a factor of degree at most four. All computations use exact arithmetic in $\mathbb{F}_p[X]$. The listed constants and the checking rule are the complete certificate; the result does not rely on a search failure or on floating-point arithmetic. $\square$

III.109.6 Proof of the classification and a quantitative consequence

Proof of Theorem III.109.1. If some $v_j = 0$, moving only that block leaves the colour constant, so avoidance fails. Suppose next that a block has even degree. Set its centre to zero and move in that block alone; norm parity gives a symmetric progression, and every other block contributes a constant. If a block has odd degree at least $2m + 1$, the same argument uses Theorem III.109.5, Theorem III.109.6, or Proposition III.109.4, as appropriate, to obtain a symmetric progression in that single block. Thus every surviving degree must be odd and smaller than $2m$.

Suppose the coefficient vectors for one such degree d are linearly dependent. Choose scalars t_j , not all zero, with $\sum_{d_j=d} v_j t_j = 0$. The field norm is surjective, so choose h_j with $N_j(h_j) = t_j$, and set the other direction blocks to zero. At centre zero,

$$C(Th) = c + T^d \sum_{d_j=d} v_j N_j(h_j) = c \quad (T \in \mathbb{F}_q).$$

This is a nontrivial monochromatic line, proving necessity of independence.

Conversely, assume both stated conditions. For a symmetric progression the vector polynomial $C(z + Th) - C(z - Th)$ has degree at most $2m - 1$ and has the $2m$ distinct roots $\pm\alpha_i$. It vanishes identically, so every odd coefficient of $C(z + Th)$ is zero. If $h \neq 0$, take the largest degree d of a moving block. Higher-degree blocks contribute constants, and the coefficient of T^d is

$$\sum_{d_j=d} v_j N_j(h_j).$$

Its vanishing and the stated linear independence force every $N_j(h_j)$ in this degree to be zero. A field norm vanishes only at zero, contradicting the choice of d . This proves avoidance. At most r independent coefficient vectors can occur at each degree, giving (III.109.2). Equality is attained by using one block of every odd degree less than $2m$ for each standard basis vector of $\mathbb{F}_q^r$, precisely r independent copies of the layered norm construction. $\square$

Corollary III.109.7 (Supersaturation survives independent additions). *Let W be an e -dimensional $\mathbb{F}_q$ -space, $R : W \rightarrow \mathbb{F}_q^r$ be arbitrary, and $v \neq 0$. For*

$$C(x, y) = vN_{\mathbb{F}_{q^d}/\mathbb{F}_q}(x) + R(y)$$

one has

$$A_{2m}^{\text{sym}}(C) \geq q^e(q^d - 1)M_{d,m}(q).$$

In every range where the preceding theorems give $M_{d,m}(q) > 0$, no independent addition can repair this offending norm block.

Proof. Hold the direction in W equal to zero. At each of its q^e centres, the reflection equalities reduce to those of the norm, because $v \neq 0$. Proposition III.109.2 counts these progressions exactly. Other directions can only increase the total. $\square$

III.109.7 Exact finite checks and the remaining problem

The companion program implements polynomial arithmetic over prime fields, verifies the degree-nine certificates by the two criteria just described, and optionally constructs entire norm tables by traversing the multiplicative group of a certified field. Modular exponentiation independently checks fifty sampled norm-table entries in each field, in addition to the forced entries. All progression counts in the following table are exhaustive; M is the ratio count in (III.109.3) and $A = (q^d - 1)M$.

q	d	$2m$	M	A
5	3	4	0	0
5	4	4	29	18096
5	5	4	220	687280
7	3	4	0	0
7	5	4	470	7898820
7	5	6	0	0
11	5	4	1600	257680000
11	5	6	0	0
7	7	4	23072	19000761024
7	7	6	3822	3147577524

The finite tables supplement the analytic proofs; only the eleven irreducibility certificates are used to prove the small-field part of the eight-term theorem.

For a fixed general m , removing the field-size condition in Theorem III.109.1 reduces to the norm of degree $2m + 1$ over primes

$$2m < p < (2m - 1)^2.$$

There are no proper prime powers in this interval with characteristic larger than $2m$. A useful next target is a uniform existence theorem for irreducible polynomials whose odd part is

$$X \prod_{i=1}^m (X^2 - (2i - 1)^2).$$

An irreducible polynomial with this odd part supplies a symmetric progression in the critical norm; a constant even part, as in the eleven displayed certificates, supplies a monochromatic one. This is a concrete finite-field question. A solution would remove the only remaining condition from the norm-block classification for all progression lengths. It would still leave the unrestricted quintic problem separate: an improvement beyond dimension four, if one exists, must use interactions outside the independent norm-block class.

Part IV

Floor patterns and the Proposition 17.7 phase extraction

Contents. The number of floor patterns of subset sums, which controls how a cube meets a partition into progressions (Corollaries 17.5–17.6 and Proposition 17.7). Sections IV.1–IV.5: source 01, delivered Sections 9–13 (exact range factorization, quadratic-exponent upper and lower bounds, $\log_2 C_k \sim k^2$ via the classical threshold-function count, affine patterns and translation classes, polynomial features, certified $C_1, \dots, C_4 = 1, 2, 10, 154$). Section IV.6: source 04, delivered Section 8 (the same counts as F_k , the stronger elementary lower bound $2^{k(k-1)/2}$, exact modular range dependence, few profiles at a fixed slope, and a local phase-extraction theorem with coefficient $2^{-k^2-k \log_2 k - O(k)}$ at the stage of Proposition 17.7).

Overlaps. Sources 01 and 04 are independent manuscripts; they share only this count ($C_k = F_k$), which they prove independently with different certificate formats. Dated notes credit the shared lemmas and record that source 01’s asymptotic answers source 04’s Question 8.

Reading conventions of this Part (the full dictionary is Section F.3).

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
C_k, F_k, A_k	01: C_k fractional carry patterns, A_k affine patterns; 04: F_k (= 01’s C_k) and A_k (= 01’s A_k).
H_k, J_k, Q_k	Both: $H_k = (k-2)2^{k-1} + 1$; 01: $J_k = k2^{k-1}$; 04: $Q_k = k2^{k-1}$ (= 01’s J_k).
region bound	01: $\mathcal{R}(h, d) = \sum_{j \leq d} \binom{h}{j}$; 04: $R_d(H) = \sum_{j \leq d} \binom{H}{j}$. <i>The arguments are in the other order: 01’s $\mathcal{R}(H, d)$ is 04’s $R_d(H)$.</i>
B_k	01: translation budget $(2(k+1))^k \mathcal{R}(k2^{k-1}, k+1)$; 04: $(2k)^k R_{k+1}(k2^{k-1})$. <i>Different numbers</i> (note after Corollary IV.6.4).
ranges, M	01: coefficients in $(-r_i, r_i)$, M a modulus, T_n threshold counts, M_n nonnegative ones; 04: (a_i, b_i) with $m_i = b_i - a_i$, M a modulus, $\Omega_k = \{0, 1\}^k$, L labels and K local multiplicity in Lemma IV.6.6.

Added in the second write (6 October 2026): sources 06–10. Contents. Section IV.7: source 06, delivered Section 2 (exact separation of range and carries for every finite $V \subset \mathbb{Z}_{\geq 0}^k$ containing 0 and the unit vectors, the modular count $\min(2r, M)^k C(V; M)$, affine carries and translation classes, and $\log_2 C_k \sim k^2$ uniformly modulo every $M \geq 2$). Section IV.8: delivered Section 4 (localized phase removal at Proposition 17.7 with coefficient $2^{-2k^2 - O(k \log k)}$). Section IV.9: delivered Section 7 (the certified carry counts $C_0, \dots, C_4 = 1, 1, 2, 10, 154$ and $D_0, \dots, D_3 = 1, 2, 6, 38$, the uniformity and density checks, and the size of the improvements).

Overlaps (dated notes). The counts C_k are certified independently by sources 01 and 06 up to $k = 4$ ($C_4 = 154$ twice, with different certificate formats); D_k agrees with 04’s A_k for $k \leq 3$. Exact range scaling (01, 04, 06), the modular count (04, 06), general coefficient families (01’s polynomial features, 06’s V), $\log_2 C_k \sim k^2$ through the threshold-function count (01, 06). At Proposition 17.7 three coefficients are proved: 04’s $2^{-k^2-k \log_2 k - O(k)}$, 06’s $2^{-2k^2 - O(k \log k)}$, and 01’s weaker selection interface (note after Corollary IV.6.8).

Reading conventions for the sections of sources 06–10 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
C_k, D_k, L_k	06: C_k ($= 01$'s $C_k = 04$'s F_k), D_k affine carry patterns ($D_1, D_2, D_3 = 2, 6, 38 = 04$'s A_1, A_2, A_3), $C_{k,M}, D_{k,M}$ their reductions modulo M , $L_k(r, M)$ translation classes, $C(V), D(V)$ for a general family V .
$h_k, \tilde{h}_k, \mathcal{R}_d(m)$	06: $h_k = (k-2)2^{k-1} + 1$ ($= H_k$ of 01, 04), $\tilde{h}_k = k2^{k-1}$ ($= 01$'s J_k , 04's Q_k), $\mathcal{R}_d(m) = \sum_{j \leq d} \binom{m}{j}$ ($= 01$'s $\mathcal{R}(m, d) = 04$'s $R_d(m)$). <i>In $\mathcal{R}_d(m)$ the subscript is the dimension.</i>
$K_k, c_k, \mathcal{E}_d$	06: $K_k = (2(k+1))^k \mathcal{R}_{k+1}(k2^{k-1})$ ($= 01$'s B_k) and $c_k = 1/(8K_k^2)$; $\mathcal{E}_d(g) = p^{d+1} \ g\ _{U^d}^{2^d}$ an unnormalized cube energy; q, M, w, ℓ the partition scales. <i>06's c_k is not 04's c_k of Corollary IV.6.8; this $\mathcal{E}_d$ is none of the $\mathcal{E}$'s of Parts II–III.</i>
$T(n), T_+(n)$	06: the number of Boolean threshold functions and of those with nonnegative weights (01: T_n, M_n).

Added in the fourth write (6 October 2026): sources 15–20. Contents. Sections IV.10–IV.14 are source 16's chapter: a guide written for this report and its delivered Sections 2–5 (cube forms and the partition quotient; explicit polarization, a factorial criterion and a functional test over $\mathbb{F}_p$; sliding-window localization without carry patterns, with its width optimization and a local-scale ceiling; a refinement that keeps the original small cells).

Overlaps (dated notes). Proposition 17.7 now has four coefficients in this Part: source 16's $(4/(3(k+2)))^{k+2}$ (the strongest at every k), source 04's, source 06's and source 01's selection interface; the comparison in one normalization is the note after Corollary IV.6.8. Source 16's Section IV.14 repeats source 04's Lemmas IV.6.5–IV.6.6 (credited at both ends); its polarization is source 04's and the formal development's. The carry-pattern counts of this Part stay needed for Corollaries 17.5–17.6.

Reading conventions for the sections of sources 15–20 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
windows	16: N prime, $m = 2h + 1$, $\mathcal{C}_j(u) = N^{j+1} \ u\ _{U^j}^{2^j}$, $\mathcal{B}_k(g; \mathcal{Q}, U)$ the partition quotient with normalization bound U (<i>not 01's or 04's B_k</i> ; U is the formal $l+1$), $c_k = (4/(3(k+2)))^{k+2}$ (<i>not 04's or 06's c_k</i>), $R = kh$, $w = N/M$, $U = \lceil w \rceil$, S_t windows, $\mathcal{P}_j$ partitions, $D(a) = \sum_i a_i $, $T = (k+2)h$, $A_d(L)$ interval cube counts, R_0 and L active and global pattern counts (04's 2^k and L), $\mathcal{C}$ conjugation.

Added in the seventh write (6 October 2026): sources 34–39. Contents. Sections IV.15–IV.24: a guide and source 36's delivered Sections 2–10 (normalization and the cube identity; the sharp bilinear estimate; the derivative rank profile; the symmetry threshold; equality and stability; lossless polarization and the Section 17 interface; robust and local variants; multiaffine selectors and extension fields; small characteristic).

Overlaps (dated notes). Source 36's bilinear estimate sharpens source 31's noncommuting-subgroup gap, and its threshold at $k = 1$ is a vector-space analogue of source 29's symmetry step; it cites neither. **Cross-source answers.** None; its questions are new (items IV.8 and IV.9).

Reading conventions for the sections of sources 34–39 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Section 17	36: $V = \mathbb{F}_p^n$, p odd, $\sigma : V^k \rightarrow V^*$ a multilinear selector, T its scalar $(k+1)$ -linear form, $\mathcal{A}_f(\sigma)$ the tested energy, $\vartheta_{k,p} = 1 - (1 - 1/p)^k$, $\mathbf{h}, \mathbf{u}$ tuples, $\mathcal{Q}_B$ a bilinear functional (<i>not $\mathbb{Q}$</i>), rad the radical (<i>not $\mathbb{R}$</i>), r half-ranks, P the canonical polynomial, $e_p(t) = e^{2\pi it/p}$.

Added in the eighth write (7 October 2026): sources 40–46. Contents. Sections IV.25–IV.41: a guide, source 46’s Sections 2–9 (conventions and the first integration obstruction; the exact binary mixed cubic norm; restriction, agreement and stability; constructive integration and phase order; all-characteristic scalar polarization; trace-multiplication tensors; restriction dimensions and repair costs; Gowers’s derivative-spectrum step) and source 43’s Sections 2–9 (cube identities and phase invariance; division-free finite-abelian orthogonality; symmetry thresholds on finite abelian groups; Boolean nonclassical integration; the shear identity and Boolean certification; exact cubic obstruction ranks; mixed correlations; the quartic frontier).

Overlaps (dated notes). Sources 43 and 46 prove the same cubic theorem and the same improvement of Tidor’s codimension independently; source 43’s orthogonality and symmetry threshold extend source 36’s to every finite abelian group. **Cross-source answers.** Source 36’s characteristic-two question is answered for symmetric trilinear forms by sources 46 and 43, and its nonclassical and composite questions are advanced by source 43 (item IV.9).

Reading conventions for the sections of sources 40–46 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Section 17	46: $V = \mathbb{F}_2^n$ (or $\mathbb{F}_p^n$), T a symmetric trilinear form, C_T , r , $m = r/2$, Λ_T the eight-function average, Γ_T the seven-function correlation, $M_T(f)$, ι the embedding $\mathbb{F}_p \rightarrow \mathbb{T}$, $T_{c,d}$ trace tensors, $\mathcal{C}$ conjugation. 43: $\mathcal{E}_T(f)$ (<i>not an additive energy</i>), Ω_T ($= 46$ ’s C_T), $\mathcal{M}(T)$, $\mathcal{Q}_B$, κ_B , R_B , d_B , U_a projective translations, nCSM^d , SymMult^d , C_d (<i>not 18’s or 44’s C_d</i>), $\mathbb{T} = \mathbb{R}/\mathbb{Z}$, $p_*(G)$.

Added in the ninth write (7 October 2026): sources 47–53. Contents. Sections IV.42–IV.50: a guide and source 47’s Sections 2–9 (finite calculus and nonclassical integration; the spectral obstruction of the defect; derivative periods and forbidden frequencies; exact canonical energies in every degree; low-complexity defects; support bounds for non-pure symmetric maps; universal thresholds; deficits and noisy selectors).

Overlaps (dated notes). Source 47 continues source 43 (the shear profile, the cubic law, the quartic reduction), which it cites by library file name. **Cross-source answers.** Source 43’s quartic question is answered (item IV.10).

Reading conventions for the sections of sources 47–53 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Section 17	47: $V = \mathbb{F}_2^n$, T a d -linear map with its evaluation slot last, $\mathcal{E}_T(f)$ the selected energy ($= 43$ ’s), $\mathcal{M}(T)$ its maximum, $\mathcal{I}_d(V)$ the integrable tensors ($= 43$ ’s nCSM^d), $C_d^{u,v}$ the canonical tensor (<i>not 18’s or 44’s C_d</i>), $b_d = 1 - (d+1)/2^d$, $M_{d,n}^{\text{sym}}$, Φ_T the frozen defect, D_{00}, D_{01}, D_1 deficits, $\mathbb{R}/\mathbb{Z} = \mathbb{R}/\mathbb{Z}$.

Added in the eleventh write (7 October 2026): sources 60–69. Contents. Sections IV.51–IV.66: a guide, source 65’s Sections 2–8 (finite calculus and integration; the cubic base case; the canonical energy; defect geometry; the induction; rigidity, repair geometry and counting; robust certification) and source 69’s Sections 2–9 (the integrable tensor space; the cubic base case; rank-one defects; the contraction induction; extremal flags; the minimal obstruction quotient; the sharp construction; energy to complexity). Sections IV.67–IV.68: a guide and source 63’s Section 6 (asymptotically sharp local phase removal). **Cross-source answers.** Source 47’s Research questions 1–2 and the tensor part of 4 are answered by sources 65 and 69 (item IV.11); source 16’s question on the best local coefficient is answered to first order by source 63 (item IV.7).

Reading conventions for the sections of sources 60–69 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
Boolean	65, 69: $V = \mathbb{F}_2^n$, $\mathcal{E}_T(f)$ the selected energy (= 47's), $b_d = 1 - (d+1)/2^d$, $C_d^{u,v}$ the canonical tensor (= 47's C_d), Φ_T the repeated-variable defect; 65: $\mathcal{M}(T)$, $\mathcal{I}_d(V)$, $\text{SymMult}^d(V)$, K_T the radical of integrable contractions, $r(T) = \text{codim } K_T$; 69: $\mathcal{M}(T)$, $\mathcal{I}_d(V)$, $\text{SM}_d(V)$, $R(T)$ (= 65's K_T), $\rho(T)$ (= 65's $r(T)$), $K(T) \subseteq R(T)$ the complete defect radical (<i>not 65's K_T</i>), $m_{\text{eff}}(T)$, $B_d(r)$.
Proposition 17.7	63: $v_k = 2^{k+1}/(k+2)!$ (= 16's ceiling), D_k , $\widehat{D}_k$ denominators, c_k^{loc} , C_r unnormalized cube sums, d the common difference (<i>not a norm order here</i>), m an odd box width, w a window width, K_w the kernel, $\mu_{k,\ell,w}$ the face measure.

IV.1 Fractional carry patterns and exact range dependence

Source 01, delivered Section 9.

We turn to the floor-pattern interface. For $k \geq 1$ and $x = (x_1, \dots, x_k) \in [0, 1)^k$, define

$$c_x(S) = \left\lfloor \sum_{i \in S} x_i \right\rfloor, \quad S \subseteq [k], \quad C_k = \#\{c_x : x \in [0, 1)^k\}. \quad (\text{IV.1.1})$$

The values at the empty set and at singletons are always zero. We call c_x a *fractional carry pattern*. The terminology refers only to this finite function; no probabilistic assumption on x is made. For completeness set $C_0 = 1$.

IV.1.1 Why boundary points do not introduce extra patterns

A hyperplane-region count normally concerns strict inequalities. Floor functions also take values on the hyperplanes, so that issue must be settled before counting chambers.

Lemma IV.1.1 (Positive perturbation). *Every fractional carry pattern has a representative $y \in (0, 1)^k$ such that $\sum_{i \in S} y_i$ is nonintegral for every nonempty $S \subseteq [k]$.*

Proof. Given $x \in [0, 1)^k$, put $y_i = x_i + \varepsilon$. There are finitely many coordinates and subset sums. Choose $\varepsilon > 0$ so small that $x_i + \varepsilon < 1$ for every i and

$$|S|\varepsilon < \left\lfloor \sum_{i \in S} x_i \right\rfloor + 1 - \sum_{i \in S} x_i \quad (\emptyset \neq S \subseteq [k]).$$

Every right-hand side is strictly positive, including when the subset sum is an integer. The perturbation moves the sum positively but not to the next integer. Its floor is unchanged, and it is no longer an integer. All coordinates become strictly positive. $\square$

Theorem IV.1.2 (Exact separation of the coefficient range). *Let $r_1, \dots, r_k$ be positive integers and let $\mathcal{F}(\mathbf{r})$ be the family of functions*

$$S \mapsto \left\lfloor \sum_{i \in S} \alpha_i \right\rfloor, \quad -r_i < \alpha_i < r_i.$$

Then

$$|\mathcal{F}(\mathbf{r})| = \left(\prod_{i=1}^k 2r_i \right) C_k. \quad (\text{IV.1.2})$$

Proof. Write $\alpha_i = m_i + x_i$, where $m_i = \lfloor \alpha_i \rfloor$ and $x_i \in [0, 1)$. Necessarily

$$m_i \in \{-r_i, -r_i + 1, \dots, r_i - 1\}.$$

For every subset,

$$F_\alpha(S) = \sum_{i \in S} m_i + c_x(S). \quad (\text{IV.1.3})$$

The singleton values of F_α recover all m_i ; after subtracting their subset sums, they also recover c_x . Thus this decomposition is injective as a coding of the resulting functions.

Conversely, take any permitted integer vector m and any carry pattern. By Lemma IV.1.1, choose a representative $x \in (0, 1)^k$. Then $-r_i < m_i + x_i < r_i$ for every i , even when $m_i = -r_i$. Equation (IV.1.3) realizes the chosen pair. Hence there is a bijection between the pattern family and the Cartesian product of the integer vectors with the C_k carry patterns. $\square$

The endpoint convention is essential for the exact statement. The proof handles the open interval $(-r_i, r_i)$ precisely because each fractional pattern admits an interior representative. For nonintegral range bounds $R_i > 0$, replacing r_i by $\lceil R_i \rceil$ gives an immediate upper bound, but generally not the same exact factorization.

IV.1.2 A hyperplane-arrangement upper bound

Inside the open unit cube, a nonempty subset S produces integer threshold hyperplanes

$$\sum_{i \in S} x_i = j, \quad 1 \leq j \leq |S| - 1.$$

There are exactly

$$H_k = \sum_{\emptyset \neq S \subseteq [k]} (|S| - 1) = k2^{k-1} - (2^k - 1) = (k - 2)2^{k-1} + 1 \quad (\text{IV.1.4})$$

such hyperplanes. Different subsets have different normal vectors, and none of these hyperplanes is a duplicate of another.

Lemma IV.1.3 (Region bound). *The number of regions cut out by h affine hyperplanes in $\mathbb{R}^d$ is at most*

$$\mathcal{R}(h, d) = \sum_{j=0}^d \binom{h}{j}.$$

The same bound holds for their regions restricted to an open convex set.

Proof. Let $R(h, d)$ be the maximal possible number of regions. When one adds a hyperplane, each newly created region comes from splitting an old region that meets it. The old hyperplanes induce on the new hyperplane an arrangement of at most $h - 1$ hyperplanes in dimension $d - 1$; coincident or empty intersections create no additional regions. Therefore

$$R(h, d) \leq R(h - 1, d) + R(h - 1, d - 1).$$

The initial values $R(0, d) = 1$ and $R(h, 0) = 1$, together with Pascal's identity, prove the bound by induction. A region is an intersection of open half-spaces and is convex. Intersecting it with an open convex set cannot split it into several nonempty connected pieces, so restriction cannot increase the number. $\square$

Theorem IV.1.4 (Quadratic-exponent carry bound). *For every $k \geq 1$,*

$$C_k \leq \mathcal{R}(H_k, k). \quad (\text{IV.1.5})$$

For $k \geq 3$ this implies

$$\log_2 C_k \leq k^2 + (\log_2 e - 1)k. \quad (\text{IV.1.6})$$

Proof. Lemma IV.1.1 realizes every carry pattern off the threshold hyperplanes. Off those hyperplanes, the floor values determine, and are determined by, the sides of all the listed integer thresholds. A given sign pattern is an intersection of open half-spaces with the cube and is convex. Thus distinct carry patterns correspond exactly to nonempty regions of this arrangement inside the open cube. Apply Lemma IV.1.3.

For $h \geq k \geq 1$, set $t = k/h \leq 1$. Then

$$t^k \sum_{j=0}^k \binom{h}{j} \leq \sum_{j=0}^k \binom{h}{j} t^j \leq (1+t)^h \leq e^k.$$

Hence $\mathcal{R}(h, k) \leq (eh/k)^k$. For $k \geq 3$, $H_k \geq k$ and $H_k \leq k2^{k-1}$. Substitute these inequalities and take binary logarithms. $\square$

Corollary IV.1.5 (A strengthened floor-pattern bound). *For positive integers r_i ,*

$$|\mathcal{F}(\mathbf{r})| \leq \left(\prod_i 2r_i \right) \mathcal{R}(H_k, k).$$

In particular, for $r_i = r$ and $k \geq 3$,

$$\log_2 |\mathcal{F}(r, \dots, r)| \leq k \log_2(2r) + k^2 + (\log_2 e - 1)k. \quad (\text{IV.1.7})$$

Proof. Combine Theorems IV.1.2 and IV.1.4. $\square$

k	H_k	old log bound, $r = 1$	new log bound, $r = 1$
2	1	16	3.000
3	5	54	7.700
4	17	128	15.650
8	769	1024	69.358
12	20481	3456	155.025
16	458753	8192	272.667

Table 5: Binary logarithms of the bounds 2^{2k^3} and $2^k \mathcal{R}(H_k, k)$ for the coefficient range $(-1, 1)^k$. These upper bounds are not asserted to be exact counts.

IV.2 Lower bounds and the optimal entropy scale

Source 01, delivered Section 10.

IV.2.1 An elementary independent-choice construction

Theorem IV.2.1 (Elementary quadratic-exponent lower bound). *For $k \geq 1$,*

$$C_k \geq 2^{\lfloor k^2/4 \rfloor}. \quad (\text{IV.2.1})$$

Proof. The assertion for $k = 1$ is immediate. Write $k = m + n$ with $m, n \geq 1$. Use m anchor coordinates

$$x_j = \frac{2^{j-1}}{2^m} \quad (1 \leq j \leq m).$$

Their subset sums are precisely $t/2^m$, for $0 \leq t \leq 2^m - 1$. For each of the remaining n coordinates independently choose

$$y_i = \frac{h_i + 1/2}{2^m}, \quad 0 \leq h_i \leq 2^m - 1.$$

For an anchor subset with sum $t/2^m$,

$$\left\lfloor y_i + \frac{t}{2^m} \right\rfloor = \begin{cases} 0, & h_i + t < 2^m, \\ 1, & h_i + t \geq 2^m. \end{cases}$$

This threshold, as t ranges over all anchor subset sums, recovers h_i : for $h_i = 0$ there is no value 1, and otherwise the first value 1 occurs at $t = 2^m - h_i$. Therefore different choices of the vector $(h_1, \dots, h_n)$ give different carry patterns. There are 2^{mn} choices. Taking $m = \lfloor k/2 \rfloor$ and $n = \lceil k/2 \rceil$ proves the theorem. $\square$

This construction is independent of any threshold-function enumeration. Together with (IV.1.6), it already determines the order $\log C_k = \Theta(k^2)$ by elementary arguments.

IV.2.2 Embedding Boolean threshold functions

Let T_n denote the number of Boolean linear threshold functions on $\{0, 1\}^n$ admitting a strict representation

$$g(\varepsilon) = \mathbf{1}_{\sum_i w_i \varepsilon_i > t}, \quad \sum_i w_i \varepsilon_i \neq t \text{ for all } \varepsilon.$$

Let M_n count the nonconstant such functions with a representation using nonnegative weights. We need only a simple relation between these two counts and carry patterns.

Lemma IV.2.2 (Thresholds inject into one higher carry dimension). *For every $n \geq 1$,*

$$C_{n+1} \geq M_n \geq 2^{-n} T_n - 2. \quad (\text{IV.2.2})$$

Proof. If a coefficient w_i is negative, replace the Boolean variable ε_i by $1 - \varepsilon_i$ and adjust the constant threshold. This makes that coefficient nonnegative. Thus every threshold function is a coordinate-flipped version of a nonnegative-weight threshold function, with at most 2^n choices of coordinate flips. Including the two constant functions gives $T_n \leq 2^n(M_n + 2)$.

For the first inequality, fix a nonconstant nonnegative-weight threshold function. Strictness and nonconstancy imply $0 < t < \sum_i w_i$. Choose $c > 0$ such that $c \sum_i w_i < 1$, and set

$$x_0 = 1 - ct, \quad x_i = cw_i \quad (1 \leq i \leq n).$$

All coordinates lie in $[0, 1)$, and $x_0 \in (0, 1)$. For $S \subseteq [n]$, the sum $x_0 + \sum_{i \in S} x_i$ lies strictly between 0 and 2, is not 1, and satisfies

$$\left\lfloor x_0 + \sum_{i \in S} x_i \right\rfloor = \mathbf{1}_{\sum_{i \in S} w_i > t}.$$

Therefore the carry pattern restricted to subsets containing coordinate 0 recovers the original Boolean function. Distinct functions have distinct resulting carry patterns, regardless of the chosen strict representations and sufficiently small scaling constants. $\square$

The classical external input is

$$\log_2 T_n = n^2 - n \log_2 n + O(n). \quad (\text{IV.2.3})$$

The leading n^2 asymptotic is due to Zuev; the refinement in (IV.2.3) is attributed to Kahn–Komlós–Szemerédi. Baldi and Vershynin explicitly record both results and also recover the leading asymptotic as the degree-one case of their theorem [8, Section 1.4 and Theorem 1.1]. We use this input, rather than claiming to reprove it; the original references are [9, 10].

Theorem IV.2.3 (Optimal leading carry entropy). *Using the classical input (IV.2.3),*

$$k^2 - k \log_2 k - O(k) \leq \log_2 C_k \leq k^2 + (\log_2 e - 1)k. \quad (\text{IV.2.4})$$

In particular $\log_2 C_k \sim k^2$. Consequently, for arbitrary positive integer ranges $r_1, \dots, r_k$,

$$\log_2 |\mathcal{F}(\mathbf{r})| = \sum_{i=1}^k \log_2(2r_i) + k^2 + o(k^2), \quad (\text{IV.2.5})$$

where the error depends only on k , not on the ranges.

Proof. For $n = k - 1$, Lemma IV.2.2 and (IV.2.3) give

$$\log_2 C_k \geq (k - 1)^2 - (k - 1) \log_2(k - 1) - O(k).$$

The subtraction of two constant functions is negligible at this scale, and the factor 2^{-n} costs only $O(k)$ in the logarithm. The right side is $k^2 - k \log_2 k - O(k)$. Combine this with Theorem IV.1.4. Finally take logarithms in the exact identity (IV.1.2); its range term has no approximation error. $\square$

The leading coefficient 1 in (IV.2.4) is therefore settled by the stated classical input. The present estimates leave a possible $k \log k$ correction unresolved. For a completely elementary version of the article, omit (IV.2.3) and replace the asymptotic claim by the bounds in (F.6.5); every local density theorem is unaffected.

[write] 6 October 2026. This theorem answers source 04’s Question 8 (dated note there, Section V.17): source 04’s F_k is C_k and its A_k is the A_k of Section IV.3. Source 04’s elementary lower bound $F_k \geq 2^{k(k-1)/2}$ (Theorem IV.6.2) is stronger than Theorem IV.2.1 for every $k \geq 3$ (note there).

[write] 6 October 2026; sources 06–10. Source 06’s Theorem IV.7.6 proves the same leading entropy through the same external threshold-function count (Zuev), and extends it: $\log_2 C_{k,M}$ and $\log_2 D_{k,M}$ are $(1 + o(1))k^2$ uniformly in the modulus $M \geq 2$, and so is the number $L_k(r, M)$ of translation classes. It does not state the secondary term $-k \log_2 k - O(k)$ given here. Credited to both.

IV.3 Affine patterns, translation classes, and general features

Source 01, delivered Section 11.

IV.3.1 Fractional affine offsets

For $t = (t_0, t_1, \dots, t_k) \in [0, 1)^{k+1}$ define

$$a_t(S) = \left\lfloor t_0 + \sum_{i \in S} t_i \right\rfloor, \quad A_k = \#\{a_t : t \in [0, 1)^{k+1}\}.$$

Here $a_t(\emptyset) = 0$. Let

$$J_k = \sum_{S \subseteq [k]} |S| = k2^{k-1} \quad (k \geq 1). \quad (\text{IV.3.1})$$

Theorem IV.3.1 (Affine fractional-pattern bound). *For $k \geq 1$,*

$$C_k \leq A_k \leq \mathcal{R}(J_k, k+1). \quad (\text{IV.3.2})$$

In particular, with the classical input of Theorem IV.2.3, $\log_2 A_k \sim k^2$.

Proof. Taking $t_0 = 0$ shows $A_k \geq C_k$. For the upper bound, perturb all $k+1$ coordinates by a common sufficiently small positive number. Each nonempty affine sum increases by $(|S|+1)\varepsilon$ but does not reach the next integer, and t_0 remains less than one. Thus every pattern has an interior representative off all integer threshold hyperplanes. For each subset S the internal thresholds are

$$t_0 + \sum_{i \in S} t_i = j, \quad 1 \leq j \leq |S|.$$

Their number is J_k . Apply Lemma IV.1.3 in dimension $k+1$. For $k \geq 2$, $J_k \geq k+1$, and the binomial-sum estimate gives

$$\log_2 A_k \leq (k+1) \log_2 \frac{ek2^{k-1}}{k+1} = k^2 + O(k).$$

Use the lower bound from C_k for the leading asymptotic. □

IV.3.2 Reduction modulo a modulus and quotienting by translations

For a positive integer M , let $\mathcal{A}_M(\mathbf{r})$ consist of

$$S \mapsto \left\lfloor \alpha_0 + \sum_{i \in S} \alpha_i \right\rfloor \pmod{M}, \quad \alpha_0 \in \mathbb{R}, \quad -r_i < \alpha_i < r_i.$$

The group $\mathbb{Z}/M\mathbb{Z}$ acts on these functions by adding a constant to all values. This action preserves the family, since α_0 is unrestricted.

Theorem IV.3.2 (Modulus-free translation-class budget). *For positive integers r_i and M ,*

$$|\mathcal{A}_M(\mathbf{r})| \leq M \left(\prod_i 2r_i \right) \mathcal{R}(J_k, k+1), \quad (\text{IV.3.3})$$

$$|\mathcal{A}_M(\mathbf{r})/(\mathbb{Z}/M\mathbb{Z})| \leq \left(\prod_i 2r_i \right) \mathcal{R}(J_k, k+1). \quad (\text{IV.3.4})$$

Proof. Write $\alpha_0 = m_0 + t_0$ and $\alpha_i = m_i + t_i$ with fractional parts in $[0, 1)$. Every pattern equals

$$m_0 + \sum_{i \in S} m_i + a_t(S) \pmod{M}.$$

There are M choices of m_0 modulo M , at most $\prod_i 2r_i$ integer coefficient vectors, and A_k fractional affine patterns. This proves (IV.3.3).

Every translation orbit has a unique representative with value zero at the empty set: subtract the original value there. This cancels m_0 while retaining the information in t_0 . The remaining possibilities are counted by at most $(\prod_i 2r_i)A_k$. Apply Theorem IV.3.1. Equivalently, the constant-translation action is free, as can be checked at the empty set, so every orbit has exactly M members. $\square$

For $M = 1$, the actual orbit count is one; the inequality remains valid. One must not model translation classes by simply discarding the fractional part of the affine offset. Canonical normalization at the empty set removes the *integer* constant modulo M , not its fractional part. Modular collisions may make either inequality strict, which is why no exact product formula is asserted in this affine modular setting.

IV.3.3 The combinatorial input for phase removal

At the relevant floor-pattern step of the proof of Proposition 17.7, the cell-index maps have affine subset-sum form, with coefficient range bounded by $k+1$ [1, proof of Proposition 17.7]. Consequently a valid translation-class budget is

$$B_k = (2(k+1))^k \mathcal{R}(k2^{k-1}, k+1). \quad (\text{IV.3.5})$$

Its logarithm is at most $k^2 + k \log_2 k + O(k)$, rather than a cubic polynomial. The following elementary statement makes precise what substituting this budget does, and what it does not do.

Proposition IV.3.3 (A transparent selection interface). *Suppose $z_1, \dots, z_L \in \mathbb{C}$ and $|\sum_{\ell=1}^L z_\ell| \geq U$. Then some ℓ satisfies $|z_\ell|^2 \geq U^2/L^2$. If the terms are indexed by translation classes from Theorem IV.3.2 with $r_i = k+1$, the denominator may be replaced by B_k^2 . More generally, any correctly established estimate $E \geq K/L^2$, with $K \geq 0$ and L the actual number of such classes, implies $E \geq K/B_k^2$.*

Proof. The triangle inequality gives $U \leq \sum_\ell |z_\ell| \leq L \max_\ell |z_\ell|$. Square, and then use $L \leq B_k$ from Theorem IV.3.2. The last assertion follows because $x \mapsto x^{-2}$ is decreasing for positive x . Zero classes cannot carry positive U ; the vacuous zero-signal case is handled separately. $\square$

[write] **6 October 2026; sources 15–20.** Source 16’s Theorem F.21.1 avoids pattern counts at this step altogether (Section IV.10).

Thus the combinatorial L^{-2} loss at this interface has a budget $2^{-2k^2-2k \log_2 k - O(k)}$. This is a local substitution conditional on the other estimates in the application being valid. We do not promote it to an improved version of every quantitative statement in Sections 17–18: localization, phase removal, uniformity normalization, and the earlier structural parameters must all be propagated separately.

IV.3.4 Nonnegative integer features and effective rank

The Boolean cube is not essential to the upper-bound mechanism. Let $V \subseteq \mathbb{N}^k \setminus \{0\}$ be a finite set of row vectors and consider

$$x \mapsto (\lfloor v \cdot x \rfloor)_{v \in V}, \quad x \in [0, 1]^k.$$

Put $H(V) = \sum_{v \in V} (\|v\|_1 - 1)$ and $d = \dim_{\mathbb{R}} \text{span } V$.

Proposition IV.3.4 (Rank-sensitive feature bound). *The number of these patterns is at most $\mathcal{R}(H(V), d)$. If V contains every coordinate unit vector, allowing $\alpha_i \in (-r_i, r_i)$ multiplies the number of fractional patterns by exactly $\prod_i 2r_i$.*

Proof. Every nonzero row has positive dot product with the all-ones vector. The common positive perturbation of Lemma IV.1.1 therefore preserves all floors while removing every integral equality and every boundary coordinate. Row v contributes at most $\|v\|_1 - 1$ internal integer hyperplanes; duplicates can only reduce the count. The linear map $x \mapsto (v \cdot x)_{v \in V}$ has rank d . Its image of the open cube is an open convex set relative to its d -dimensional image, and the relevant threshold hyperplanes restrict to affine hyperplanes in that image. Lemma IV.1.3 proves the bound. For the last assertion, $v \cdot m$ is an integer and the unit-vector coordinates recover m_i . Interior representatives prove surjectivity exactly as in Theorem IV.1.2. $\square$

Example IV.3.5 (Why arbitrary signed features need a boundary argument). Near $x = 0$, the map $x \mapsto (\lfloor x \rfloor, \lfloor -x \rfloor)$ has three patterns: $(0, -1)$ for $x > 0$, $(-1, 0)$ for $x < 0$, and $(0, 0)$ at zero. The one hyperplane has only two open regions. Thus a chamber count alone misses a floor pattern when opposite orientations occur. A sufficient replacement for nonnegative rows is a common direction u with $v \cdot u > 0$ for every row, together with room to perturb the parameter domain in that direction. Without such a condition one must count the relevant boundary strata as well.

[write] **6 October 2026; sources 06–10.** Source 06's Theorem IV.7.3 is the same separation of range and carries for any finite $V \subset \mathbb{Z}_{\geq 0}^k$ that contains 0 and the unit vectors, with the modular count $\min(2r, M)^k C(V; M)$ and the region bound $\mathcal{R}_k(h(V))$ in the ambient dimension k ; Proposition IV.3.4 is rank-sensitive (the dimension d of the feature span). Both sources make the same point about signed families (the example above; the paragraph after source 06's theorem).

IV.4 A polynomial generalization on the Boolean cube

Source 01, delivered Section 12.

The feature bound also gives an extension in which the input dimension and coefficient dimension are different. It is useful when a localization step produces low-degree polynomial floor functions rather than affine ones.

Fix $1 \leq d \leq n$ and a positive integer r . Let $\mathcal{P}_{n,d}(r)$ be the family of integer-valued functions on $\{0, 1\}^n$ of the form

$$\varepsilon \mapsto \left| \sum_{\substack{S \subseteq [n] \\ |S| \leq d}} \alpha_S \prod_{i \in S} \varepsilon_i \right|, \quad -r < \alpha_S < r. \quad (\text{IV.4.1})$$

The constant monomial corresponds to $S = \emptyset$. Put

$$m = \sum_{j=0}^d \binom{n}{j}, \quad H_{n,d} = \sum_{j=1}^d \binom{n}{j} 2^{n-j}.$$

Let $T(n, d)$ count strict Boolean polynomial threshold functions of degree at most d ; on a Boolean cube a polynomial can be multilinearized without changing its values or increasing its degree.

Theorem IV.4.1 (Polynomial floor-pattern entropy). *For $1 \leq d \leq n$ and positive integer r ,*

$$T(n, d) \leq |\mathcal{P}_{n,d}(r)| \leq (2r)^m \mathcal{R}(H_{n,d}, m). \quad (\text{IV.4.2})$$

For fixed $d \geq 1$, using the polynomial threshold-function theorem of Baldi–Vershynin, if $r = r(n)$ is a positive integer with $\log r = o(n)$, then

$$\log_2 |\mathcal{P}_{n,d}(r)| = (1 + o(1)) n \sum_{j=0}^d \binom{n}{j} = (1 + o(1)) \frac{n^{d+1}}{d!}. \quad (\text{IV.4.3})$$

Proof. Index coefficient space by the m monomials of degree at most d . For an input ε , the evaluation vector has entries $\prod_{i \in S} \varepsilon_i \in \{0, 1\}$. It is nonzero because of the constant monomial. The sum of its coordinates counts the active monomials at ε . Summing the number of internal thresholds over all inputs gives

$$\sum_{\varepsilon \in \{0,1\}^n} (\#\{S : |S| \leq d, S \subseteq \text{supp } \varepsilon\} - 1) = \sum_{j=1}^d \binom{n}{j} 2^{n-j} = H_{n,d}.$$

Separate each coefficient into an integer part and a fractional part. There are at most $(2r)^m$ integer vectors. For each vector, fractional patterns are bounded by the hyperplane count in dimension m , using positive perturbation and Lemma IV.1.3. This proves the upper bound. It does not assert that different integer vectors always give disjoint pattern families.

For the lower bound, take a strict multilinear polynomial threshold function, represented by a polynomial p with no zero values on the Boolean cube. Multiply p by a sufficiently small positive c so that every coefficient of cp lies in $(-r, r)$ and $0 < |cp(\varepsilon)| < 1$ at every input. Then

$$\lfloor cp(\varepsilon) \rfloor = \begin{cases} -1, & p(\varepsilon) < 0, \\ 0, & p(\varepsilon) > 0. \end{cases}$$

Thus the integer-valued floor function encodes the threshold function. Distinct threshold functions give distinct floors.

For fixed d , the external theorem gives $\log_2 T(n, d) = (1 + o(1))nm$ [8, Theorem 1.1]. For sufficiently large n , $H_{n,d} \geq m$ and $H_{n,d} \leq 2^n m$. The binomial estimate from Theorem IV.1.4 therefore yields

$$\log_2 |\mathcal{P}_{n,d}(r)| \leq m \log_2(2r) + m \log_2(eH_{n,d}/m) \leq nm + m \log_2(2r) + m \log_2 e.$$

Since $\log r = o(n)$, this matches the lower bound to leading order. Finally $m = (1 + o(1))n^d/d!$ for fixed d . $\square$

The finite upper bound in (IV.4.2) is elementary; only the matching leading lower asymptotic uses the external threshold theorem. This generalization also shows why using the square of the coefficient dimension as a universal exponent can be wasteful: for $d > 1$, the sharp leading entropy is on the scale n^{d+1} rather than $m^2 \asymp n^{2d}$.

IV.5 Exact finite enumeration and reproducible checks

Source 01, delivered Section 13.

IV.5.1 Small carry counts

Proposition IV.5.1 (Certified small dimensions). *The first four carry counts are*

$$C_1 = 1, \quad C_2 = 2, \quad C_3 = 10, \quad C_4 = 154. \quad (\text{IV.5.1})$$

The first three values have the elementary proof below. The fourth is proved by an exhaustive finite rational-certificate enumeration supplied with this article; its certificate format and exact verification argument are given in the following subsection.

Proof of the first three values. For $k = 1$ every floor is zero. For $k = 2$, only $\lfloor x_1 + x_2 \rfloor \in \{0, 1\}$ varies, and both values occur.

For $k = 3$, Lemma IV.1.1 allows us to avoid integral sums. If $x_1 + x_2 + x_3 < 1$, every pair floor is zero. If the total is greater than 2, every pair sum is greater than 1, since the remaining coordinate is less than 1. These give two patterns. When the total lies in $(1, 2)$, all eight assignments of pair floors occur. Up to coordinate permutations, representatives with respectively zero, one, two, and three pair floors equal to one are

$$(0.4, 0.4, 0.4), \quad (0.7, 0.5, 0.1), \quad (0.8, 0.4, 0.4), \quad (0.6, 0.6, 0.6).$$

There are $1 + 3 + 3 + 1 = 8$ such pair patterns, all with total floor one. The two outer patterns are realized by $(0.1, 0.1, 0.1)$ and $(0.9, 0.9, 0.9)$. Hence $C_3 = 10$. $\square$

For example, Theorem IV.1.2 now gives exactly $C_4(2r)^4 = 154(2r)^4$ patterns with four coefficients in $(-r, r)$. At $r = 1$ the exact count is 2464, while the simple new arrangement upper bound is 51424. No assertion is made that these finite values or their sequence have not appeared elsewhere.

[write] 6 October 2026. $C_1, C_2, C_3 = 1, 2, 10$ are certified independently in source 04 (as F_1, F_2, F_3 , Section IV.6), together with $A_1, A_2, A_3 = 2, 6, 38$. $C_4 = 154$ is certified only here. At intake the sequence 1, 2, 10, 154 matched only OEIS A134088 (1, 1, 2, 10, 154, 7994, $\dots$, a power-series coefficient sequence); its next term 7994 differs from the uncertified recount $C_5 = 8410$ (note in Section IV.6), so the match looks like a coincidence. 2, 6, 38 gave no relevant hit. This is the intake’s search of 6 October 2026, not repeated at the write; nothing was submitted to the OEIS.

[write] 6 October 2026; sources 06–10. *Stale since the second write: $C_4 = 154$ is also certified by source 06.* The sentence “ $C_4 = 154$ is certified only here” was true of sources 01–05. Source 06 certifies $C_0, \dots, C_4 = 1, 1, 2, 10, 154$ and $D_0, \dots, D_3 = 1, 2, 6, 38$ with its own exact format (215 rational chamber witnesses and 1,165 nonnegative dual certificates, checked by a standard-library verifier; Section IV.9), so $C_4 = 154$ is now certified twice, independently. Source 06’s D_k (affine carries $\lfloor t_0 + v \cdot t \rfloor$) is the A_k here and in source 04 for $k \leq 3$ by the values; source 06 does not compute $D_4 = A_4$ or C_5 , which remain uncertified (Section V.2, item IV.2).

IV.5.2 A solver-independent rational certificate

Order the H_k hyperplanes by increasing binary mask of S , and then by increasing integer threshold. Build a binary tree of strict signs. For each partial sign word, use variables $z = (x_1, \dots, x_k, \epsilon)$ and

inequalities

$$\begin{array}{ll} -x_i + \epsilon \leq 0, & x_i + \epsilon \leq 1, \\ v \cdot x + \epsilon \leq j & \text{for the sign } v \cdot x < j, \\ -v \cdot x + \epsilon \leq -j & \text{for the sign } v \cdot x > j. \end{array}$$

Write this finite system as $Az \leq b$. A sign word is strictly feasible exactly when this system has a solution with $\epsilon > 0$.

There are two certificate types. A *primal certificate* is a rational vector z satisfying $Az \leq b$ and $z_{k+1} > 0$. A *dual certificate* is a rational vector $\lambda \geq 0$ satisfying

$$\lambda^T A = (0, \dots, 0, 1), \quad \lambda^T b \leq 0. \quad (\text{IV.5.2})$$

The latter proves that every solution of $Az \leq b$ has $\epsilon = \lambda^T Az \leq \lambda^T b \leq 0$, so the sign word is strictly infeasible. This implication needs no floating-point tolerance and no appeal to an optimizer’s status message.

The generator uses a numerical linear program only to propose rational primal or dual vectors. Every proposal is checked with exact rational arithmetic before being accepted. It branches both ways at every feasible partial word and never branches below an infeasible word. The verifier reconstructs all inequalities from the sign words, checks every rational identity, checks that both children of every feasible nonterminal node are present, and rejects extra or unreachable nodes. It then counts feasible leaves. By induction on tree depth, this is an exhaustive count of all regions in the open cube. Theorem IV.1.4’s region-pattern correspondence turns that count into C_k .

Dimension k	Hyperplanes H_k	Certificate records	Feasible leaves C_k
1	0	1	1
2	1	3	2
3	5	49	10
4	17	2061	154

Table 6: Exact certificate results in `certificates/carry.json`. The count of records includes both primal and dual certificates.

IV.5.3 Commands and the limits of computational evidence

From the package directory, verification uses only the Python standard library:

```
python code/carry_certificates.py verify certificates/carry.json
```

Regeneration additionally needs NumPy and SciPy:

```
python code/carry_certificates.py generate --max-k 4 \
--out certificates/carry.json
```

The supplementary checks use NumPy and SymPy:

```
python code/check_refinements.py
```

They verify six symbolic identities, 210 broken-chord parameter cases, 210 equality constructions, 175 maximum-density root cases, 2000 sampled weighted-cell inequalities, and 24 concrete progression-partition cases. The parameter grids and random seed are fixed in the script. These are useful error detectors, not substitutes for the general proofs. In contrast, the finite carry enumeration has exact certificates for every branch and a mathematical completeness argument, though its implementation is still ordinary Python rather than kernel-checked proof code.

The verifier is independent of the numerical optimizer, not an independent second implementation of every aspect of the encoding. Its small trusted computational core includes Python integer and rational arithmetic, the constraint generator, and the tree checks. This makes the remaining trust boundary explicit.

IV.6 Floor patterns and efficient extraction of a progression profile

Source 04, delivered Section 8.

Corollaries 17.5 and 17.6 of Gowers’s paper bound the number of functions obtained by taking integer parts of linear expressions on a Boolean cube [1]. These estimates control the number of ways that a cube can meet a partition into arithmetic progressions. We isolate the fractional parts of the coefficients, determine the exact dependence on their integer ranges, and prove matching quadratic orders of magnitude for the logarithm of the remaining complexity. A second observation reduces the loss incurred when selecting one progression profile.

The edited repository already uses a joint hyperplane count in its proof of Proposition 17.7 [23]. The refinements below go further by separating integer and fractional parameters and distinguishing the number of globally possible profiles from the number possible at a fixed slope. The arguments establish the stated bounds; they do not establish historical priority for these elementary refinements.

IV.6.1 Fractional patterns and their boundary values

Write $\Omega_k = \{0, 1\}^k$ and put

$$\begin{aligned} F_k &= \# \left\{ (\lfloor \epsilon \cdot t \rfloor)_{\epsilon \in \Omega_k} : t \in [0, 1)^k \right\}, \\ A_k &= \# \left\{ (\lfloor t_0 + \epsilon \cdot t \rfloor)_{\epsilon \in \Omega_k} : (t_0, t) \in [0, 1)^{k+1} \right\}. \end{aligned}$$

Set $F_0 = A_0 = 1$, and define

$$R_d(H) = \sum_{j=0}^d \binom{H}{j}, \quad H_k = (k-2)2^{k-1} + 1, \quad Q_k = k2^{k-1},$$

where $\binom{H}{j} = 0$ for $j > H$.

[write] **6 October 2026.** *Source defect corrected in the write.* In both rows of the display above (first wording “In both displays above”, corrected after the independent check of 6 October 2026) the delivered source writes the subscript `\epsilon\in\Omega_k`, without the backslash before `epsilon`, so the delivered PDF prints the word “epsilon”; the intended $\epsilon \in \Omega_k$ is restored here. Notation: F_k is source 01’s C_k , A_k is source 01’s A_k , Q_k is source 01’s J_k , and $R_d(H)$ is source 01’s $\mathcal{R}(H, d)$, with the arguments in the other order.

Lemma IV.6.1 (Boundary removal). *Every pattern counted by F_k or A_k has a representative in the open parameter cube which lies on none of its nonconstant integer-threshold hyperplanes.*

Proof. Add a sufficiently small positive multiple of the all-ones vector to the parameters. Every nonconstant form under consideration increases strictly. An integer value moves into the interior of the same floor interval; a noninteger value remains in its floor interval if the step is sufficiently small. There are finitely many forms and every coordinate is initially less than 1, so one positive step works for all forms and coordinates. The identically zero form in the definition of F_k requires no adjustment. $\square$

Theorem IV.6.2 (Quadratic order of fractional complexity). *For every $k \geq 1$,*

$$2^{k(k-1)/2} \leq F_k \leq R_k(H_k), \quad \prod_{j=0}^{k-1} (2^j + 1) \leq A_k \leq R_{k+1}(Q_k). \quad (\text{IV.6.1})$$

Consequently $\log_2 F_k = \Theta(k^2)$ and $\log_2 A_k = \Theta(k^2)$. In particular,

$$F_k \leq e^k 2^{k(k-1)} \quad (k \geq 3), \quad A_k \leq e^{k+1} 2^{k^2-1} \quad (k \geq 2). \quad (\text{IV.6.2})$$

Proof. For a nonempty subset $S \subseteq \{1, \dots, k\}$ of size s , the form $\sum_{i \in S} t_i$ crosses precisely the possible thresholds $1, \dots, s-1$ inside the open cube. The resulting arrangement has

$$\sum_{s=1}^k \binom{k}{s} (s-1) = k2^{k-1} - (2^k - 1) = H_k$$

hyperplanes. A pattern is constant in each chamber. Its strict feasible region is an intersection of open halfspaces and is therefore convex, so a single pattern cannot occupy two chambers. Lemma IV.6.1 also accounts for all patterns originally attained on boundaries. Thus F_k is the number of chambers meeting the open cube. The standard arrangement bound, also stated as Lemma 17.4 in [1], gives $F_k \leq R_k(H_k)$.

With an intercept, the form $t_0 + \sum_{i \in S} t_i$ has thresholds $1, \dots, s$, and there are $Q_k = \sum_s s \binom{k}{s}$ relevant hyperplanes in dimension $k+1$. This gives the bound for A_k . For $H \geq d \geq 1$, $R_d(H) \leq (eH/d)^d$: apply $(1+u)^H \leq e^{Hu}$ with $u = d/H$, observing that $u^j \geq u^d$ for $j \leq d$. Applying this with $H_k/k \leq 2^{k-1}$ and $Q_k/(k+1) \leq 2^{k-1}$ proves (IV.6.2).

For the lower bound, fix one pattern counted by F_{k-1} . Its region contains a nonempty open set. Choose a representative in that set whose 2^{k-1} subset sums have pairwise distinct fractional parts. Such a choice is possible because the prohibited equalities are finitely many proper hyperplanes

$$\sum_i (\epsilon_i - \epsilon'_i) t_i = j, \quad \epsilon \neq \epsilon', \quad j \in \mathbb{Z}.$$

Only finitely many j can occur in the bounded parameter cube. In particular, no nonempty subset sum has fractional part zero. As t_k varies through $(0, 1)$, each nonempty old subset sum produces one jump, at $1 - \{\sum_{i \in S} t_i\}$. The $2^{k-1} - 1$ jump locations are distinct, and hence there are 2^{k-1} different extension patterns. Extensions of different old patterns remain different on the face $\epsilon_k = 0$. Therefore $F_k \geq 2^{k-1} F_{k-1}$, which proves the first lower bound.

For A_k , choose an old representative for which the 2^{k-1} affine sums $t_0 + \sum_{i \in S} t_i$ have distinct nonzero fractional parts. Varying t_k now gives 2^{k-1} distinct jumps and $2^{k-1} + 1$ extensions. Thus $A_k \geq (2^{k-1} + 1) A_{k-1}$, proving the second lower bound. $\square$

[write] **6 October 2026.** *Credits and a comparison.* The upper bound $F_k \leq R_k(H_k)$ is source 01's Theorem IV.1.4, and $F_k \leq e^k 2^{k(k-1)}$ is the same as 01's $\log_2 C_k \leq k^2 + (\log_2 e - 1)k$; the affine upper bound is 01's Theorem IV.3.1. The lower bound $2^{k(k-1)/2}$ is stronger than source 01's elementary $2^{\lfloor k^2/4 \rfloor}$ (Theorem IV.2.1) for every $k \geq 3$, and equal to it for $k = 1, 2$. *Proof:* $k(k-1)/2 - k^2/4 = k(k-2)/4$, which is positive for $k \geq 3$, and $\lfloor k^2/4 \rfloor \leq k^2/4$; for $k = 1, 2$ both exponents are 0 and 1. The product lower bound for A_k is source 04's alone. Using the classical threshold-function count, source 01 proves $\log_2 F_k \sim k^2$ and $\log_2 A_k \sim k^2$ (Theorems IV.2.3 and IV.3.1). This answers Question 8 below (dated note there).

IV.6.2 Exact dependence on integer ranges

Theorem IV.6.3 (Integer-part scaling). *Let $a_i < b_i$ be integers and $m_i = b_i - a_i$. The number of functions*

$$\epsilon \mapsto \left\lfloor \sum_{i=1}^k \epsilon_i \alpha_i \right\rfloor, \quad \alpha_i \in (a_i, b_i),$$

is exactly $F_k \prod_i m_i$. In particular, the count in Gowers's Corollary 17.5 is exactly $(2r)^k F_k$ for a positive integer r .

Let $F_{k,M}$ count the fractional patterns in the definition of F_k after reduction modulo a positive integer M . The corresponding number of functions with values in $\mathbb{Z}/M\mathbb{Z}$ is exactly

$$F_{k,M} \prod_{i=1}^k \min(m_i, M). \tag{IV.6.3}$$

Moreover, $F_{k,M} = F_k$ whenever $M \geq k$.

Proof. Write $\alpha_i = n_i + t_i$, where $n_i = \lfloor \alpha_i \rfloor$ and $0 \leq t_i < 1$. Then

$$\lfloor \epsilon \cdot \alpha \rfloor = \epsilon \cdot n + \lfloor \epsilon \cdot t \rfloor.$$

The singleton values recover n_i exactly, since $\lfloor t_i \rfloor = 0$. Subtracting their linear contribution recovers the fractional pattern. Conversely, every fractional pattern has an interior representative by Lemma IV.6.1; this representative works in every integer cell, including those with an excluded left endpoint. There are precisely $\prod_i m_i$ integer choices.

Modulo M , the same singleton argument recovers the residues of the n_i , and each interval of m_i consecutive integers has $\min(m_i, M)$ residues. Subtracting these residues recovers the fractional residue pattern, proving (IV.6.3). Finally, every fractional floor belongs to $\{0, \dots, k-1\}$; reduction modulo $M \geq k$ is injective on this range. $\square$

For arbitrary real endpoints, the same argument gives the upper bound $F_k \prod_i (\lceil b_i \rceil - \lfloor a_i \rfloor)$; equality need not hold because the extreme integer cells may be truncated.

[write] **6 October 2026.** With $a_i = -r_i$, $b_i = r_i$ the first assertion is source 01's Theorem IV.1.2, by the same decomposition and interior-representative argument (source 01's Lemma IV.1.1 is Lemma IV.6.1); the arbitrary integer intervals and the modular count (IV.6.3) are source 04's.

Corollary IV.6.4 (Affine modular bound). *With the integer endpoints of Theorem IV.6.3, let α_0 vary over all real numbers. The number of functions*

$$\epsilon \mapsto \lfloor \alpha_0 + \epsilon \cdot \alpha \rfloor \pmod{M}$$

is at most $MA_k \prod_i \min(m_i, M)$. The number of equivalence classes under addition of a common constant is at most $A_k \prod_i \min(m_i, M)$. Both assertions remain valid if every slope interval (a_i, b_i) is replaced by $[a_i, b_i]$.

Proof. Writing $\alpha_0 = n_0 + t_0$ gives

$$\lfloor \alpha_0 + \epsilon \cdot \alpha \rfloor = n_0 + \epsilon \cdot n + \lfloor t_0 + \epsilon \cdot t \rfloor.$$

Choose $n_0 \pmod{M}$, the allowed slope residues, and one of the A_k fractional patterns. This proves the upper bound. Translation by a constant acts freely, as seen at $\epsilon = 0$, and all M translates occur by changing n_0 .

To handle closed slope intervals, move each endpoint slope inward with velocity $d_i = 1$ at a_i and $d_i = -1$ at b_i , leaving other slopes fixed. Move t_0 upward with velocity $k + 1$. Every affine form has strictly positive velocity $k + 1 + \epsilon \cdot d \geq 1$. A sufficiently small movement preserves every floor and keeps $t_0 < 1$, while moving all slope endpoints into the open intervals. $\square$

In the symmetric case, Corollary IV.6.4 replaces the bound $M2^{2r(k+1)^3}$ in Gowers's Corollary 17.6 by

$$M \min(2r, M)^k R_{k+1}(k2^{k-1}).$$

The dependence on r is polynomial, and the exponent governing the fractional complexity is quadratic in k . Theorem IV.6.2 shows that this quadratic order cannot be reduced to $o(k^2)$.

[write] **6 October 2026.** Source 01's Theorem IV.3.2 gives the two bounds of Corollary IV.6.4 with $\prod_i 2r_i$ and $\mathcal{R}(k2^{k-1}, k + 1)$ in place of $\prod_i \min(m_i, M)$ and A_k ; source 04's are never larger. *Notation clash* B_k . Source 01's translation budget $B_k = (2(k + 1))^k \mathcal{R}(k2^{k-1}, k + 1)$ (IV.3.5) and source 04's $B_k = (2k)^k R_{k+1}(k2^{k-1})$ (Theorem IV.6.7) differ. Source 01 uses Gowers's coefficient range $k + 1$ with open intervals; source 04 uses the closed slope interval $[-k, k]$ of its phase-extraction proof, through the closed-interval clause of this corollary. Both are valid under their hypotheses. Source 04's is smaller by the factor $((k + 1)/k)^k < e$.

IV.6.3 Small cases with exact certificates

The first values are

k	1	2	3
F_k	1	2	10
A_k	2	6	38

For example, in the definition of F_3 the three pair-sum threshold indicators attain all eight Boolean patterns. If all three are zero, the total floor can be 0 or 1; if all are one, it can be 1 or 2. For each of the other six patterns, the total floor is 1. This gives $F_3 = 2 + 2 + 6 = 10$. Similarly, the two singleton values for A_2 give two possibilities for the pair value when they agree and one when they disagree, giving $A_2 = 6$.

The accompanying files `pattern_census.py` and `pattern_census.json` certify all six counts. The arrangement is constructed one hyperplane at a time, and both possible children of every surviving chamber are checked. For a system of strict inequalities $Cx < b$, a feasible branch has a rational point satisfying all inequalities strictly. An infeasible branch has rational weights $y \geq 0$ satisfying

$$\sum_i y_i = 1, \quad C^\top y = 0, \quad b^\top y \leq 0.$$

Multiplying the strict inequalities by these weights would yield $0 < b^\top y \leq 0$, a contradiction. All 450 branch certificates are checked with exact fractions. Numerical linear programming proposes the certificates during generation; the verifier itself uses only Python’s standard library and runs with

```
python -S pattern_census.py --verify pattern_census.json
```

[write] **6 October 2026.** Source 01 certifies the same F_1, F_2, F_3 (its C_k) and also $C_4 = F_4 = 154$ (Proposition IV.5.1), with a different certificate format. The census is shipped as `code/04-restriction-cubes-pattern_census.py` with `data/04-restriction-cubes-pattern_census.json`; the command above runs on a copy with the delivered names (README). At intake an independent recount reproduced $C_1, \dots, C_4$ and A_1, A_2, A_3 , and found $C_5 = 8410$ and $A_4 = 770$. It used value-branching floating-point linear programs and re-evaluated every feasibility witness exactly, but the infeasibility side rests on floating-point solver status, so these two values are *not certified*; they are recorded in Section V.2, item IV.2. They are consistent with the bounds: $2^{10} = 1024 \leq C_5 \leq R_5(H_5) = R_5(49) = 2138410$ and $2 \cdot 3 \cdot 5 \cdot 9 = 270 \leq A_4 \leq R_5(Q_4) = R_5(32) = 242825$.

[write] **6 October 2026; sources 06–10.** Source 06 (Section IV.9) certifies $C_4 = 154$ independently of source 01, and $D_1, D_2, D_3 = 2, 6, 38$ (its affine carry counts, the A_k here); it does not compute C_5 or $D_4 = A_4$, so the two intake values stay uncertified.

IV.6.4 Few profiles at a fixed slope

Lemma IV.6.5 (Intercept profiles). *For fixed $c \in \mathbb{R}^k$, the functions*

$$\epsilon \mapsto \lfloor \beta + \epsilon \cdot c \rfloor, \quad 0 \leq \beta < 1,$$

take at most 2^k different values. The same bound holds after reduction modulo M . Over the integers the count is exactly one plus the number of distinct nonzero fractional parts among the nonempty subset sums $\epsilon \cdot c$.

Proof. Each noninteger subset sum gives one jump in $(0, 1)$, at $1 - \{\epsilon \cdot c\}$. Integer subset sums give no interior jump. There are at most $2^k - 1$ distinct jump locations. The floor function is right-continuous, so a jump does not create an extra isolated pattern. Reduction modulo M can only identify patterns. Generic slopes make all subset sums distinct modulo 1, showing that the integer bound is sharp. $\square$

Lemma IV.6.6 (Extraction from locally few classes). *Let $z_{s,a}$ be complex numbers on finite index sets, with each pair (s, a) assigned one of L labels. Suppose that at most K labels occur for each fixed a . Set*

$$E = \sum_a \left| \sum_s z_{s,a} \right|^2, \quad E_j = \sum_a \left| \sum_{s: j(s,a)=j} z_{s,a} \right|^2.$$

Then $E \leq K \sum_j E_j$, and some label satisfies $E_j \geq E/(KL)$.

Proof. For each a , apply Cauchy–Schwarz to its at most K class sums. Sum over a , then average over the L labels. $\square$

[write] **6 October 2026; sources 15–20.** Source 16’s Lemma IV.14.1 and Proposition IV.14.2 reprove Lemma IV.6.5, this lemma and the first inequality of Theorem IV.6.7, without credit (note there).

IV.6.5 A quantitative local phase-extraction statement

Here is a precise form of the improvement at the stage of Proposition 17.7. For a function u on $G = \mathbb{F}_p$, write

$$\mathcal{U}_s(u) = \sum_{x, h_1, \dots, h_s \in G} \prod_{\omega \in \Omega_s} \mathcal{C}^{|\omega|} u(x - \omega \cdot h),$$

where $\mathcal{C}$ denotes complex conjugation. This quantity is nonnegative; it is $p^{s+1} \|u\|_{U^s}^2$. We use the usual Gowers–Cauchy–Schwarz inequality for its mixed version [1, Lemma 3.8].

Theorem IV.6.7 (Local extraction with a quadratic exponential loss). *Let $k \geq 1$, let $p > k + 1$ be prime, and let $m = 2h + 1 > 3k$ satisfy $m \leq \sqrt{p}$. Put $I = \{-h, \dots, h\} \subseteq G$. Let $r \in G^k$, let $\tau : G^k \rightarrow G$ be a multi-affine polynomial, and let $f : G \rightarrow \mathbb{C}$. For $\epsilon \in \Omega_k$, put $f_\epsilon(s) = f(s - \epsilon \cdot r)$. Suppose $\alpha > 0$ and*

$$\sum_{a \in I^k} \left| \sum_{s \in G} e_p(-s\tau(a)) \prod_{\epsilon \in \Omega_k} \mathcal{C}^{|\epsilon|} f_\epsilon(s - \epsilon \cdot a) \right|^2 \geq \alpha p^2 m^k, \quad e_p(t) = \exp(2\pi i t/p). \quad (\text{IV.6.4})$$

Set

$$q = \lceil h/k \rceil, \quad M = \lfloor p/q \rfloor, \quad B_k = (2k)^k R_{k+1} (k2^{k-1}).$$

There is a polynomial P of degree at most $k + 1$ and a partition $(R_i)_{i \in \mathbb{Z}/M\mathbb{Z}}$ into modular arithmetic progressions of lengths q or $q + 1$, each at least $m/(3k)$, such that

$$\sum_i \mathcal{U}_{k+1}(1_{R_i} f e_p(-P)) \geq \frac{\alpha p^2 m^k}{2^k B_k M} \geq \frac{\alpha}{9k^2 2^k B_k} m^{k+2} M. \quad (\text{IV.6.5})$$

Proof. Put $w = p/M$. Since $q \leq \sqrt{p}$, division $p = qM + r_0$ gives $0 \leq r_0 < q \leq M$. The intervals

$$Q_i = \{x \in \{0, \dots, p-1\} : iw \leq x < (i+1)w\}, \quad 0 \leq i < M,$$

therefore form a partition with lengths q or $q + 1$. Also $w \geq q \geq h/k \geq m/(3k)$, where the last inequality follows from $h = (m-1)/2$ and $m \geq 3$.

For each (s, a) , the unique progression profile of its cube is

$$j_{s,a}(\epsilon) = \lfloor s/w - \epsilon \cdot a/w \rfloor \pmod{M}.$$

This formula is valid across modular wraparound because $p/w = M$. Normalize by subtracting $j_{s,a}(0)$. The resulting profile has intercept $\{s/w\}$ and slopes $-a_i/w \in [-k, k]$. Corollary IV.6.4 gives at most $L \leq B_k$ normalized profiles globally. Lemma IV.6.5 gives at most 2^k for each fixed a .

Apply Lemma IV.6.6 to the integrand in (IV.6.4). There is a normalized profile $j : \Omega_k \rightarrow \mathbb{Z}/M\mathbb{Z}$ whose class sum has squared energy at least $\alpha p^2 m^k / (2^k L)$. Write this class sum as a sum over its M translates. A further Cauchy–Schwarz inequality gives

$$\sum_i \sum_{a \in I^k} |T_i(a)|^2 \geq \frac{\alpha p^2 m^k}{2^k L M}, \quad (\text{IV.6.6})$$

where

$$T_i(a) = \sum_s e_p(-s\tau(a)) \prod_{\epsilon \in \Omega_k} \mathcal{C}^{|\epsilon|} (1_{Q_{j(\epsilon)+i}} f_\epsilon)(s - \epsilon \cdot a).$$

The identity of this expression with the class sum uses the disjointness of the Q_i : exactly one product of their indicators is nonzero for any given (s, a) .

Extend the sum in (IV.6.6) to all $a \in G^k$, which increases it. We record the phase decomposition explicitly. For suitable polynomials P_ω of degree at most $k + 1$,

$$u\tau(a) = \sum_{\omega \in \Omega_{k+1}} (-1)^{|\omega|} P_\omega(s - \omega \cdot (a, u)). \quad (\text{IV.6.7})$$

Indeed, expand τ into squarefree monomials. For a monomial involving the coordinates indexed by $J \subseteq \{1, \dots, k + 1\}$, use

$$\prod_{j \in J} b_j = \frac{1}{|J|!} \sum_{\nu \in \{0,1\}^J} (-1)^{|\nu|} \left(s - \sum_{j \in J} \nu_j b_j \right)^{|J|}.$$

Here J includes the coordinate of u ; all required factorials are invertible because $p > k + 1$. Adding these identities proves (IV.6.7).

Expand $|T_i(a)|^2$ and write its two base points as s and $s - u$. Using (IV.6.7) makes the resulting expression the mixed $(k + 1)$ -cube form of

$$h_{i,\omega} = 1_{Q_{j(\epsilon)+i}} f_\epsilon e_p(-P_\omega), \quad \omega = (\epsilon, b), \quad b \in \{0, 1\}.$$

Gowers–Cauchy–Schwarz and the arithmetic–geometric mean inequality give

$$\sum_{a \in G^k} |T_i(a)|^2 \leq \prod_{\omega} \mathcal{U}_{k+1}(h_{i,\omega})^{1/2^{k+1}} \leq 2^{-(k+1)} \sum_{\omega} \mathcal{U}_{k+1}(h_{i,\omega}).$$

For a fixed ω , summing over i permutes all blocks, so that $\sum_i \mathcal{U}_{k+1}(h_{i,\omega})$ is independent of $j(\epsilon)$. Averaging over ω and using (IV.6.6) produces one $\omega = (\epsilon, b)$ with

$$\sum_i \mathcal{U}_{k+1}(1_{Q_i} f_\epsilon e_p(-P_\omega)) \geq \frac{\alpha p^2 m^k}{2^k L M}.$$

Finally put $t = \epsilon \cdot r$, replace $P_\omega(z)$ by $P_\omega(z + t)$, and translate every block by $-t$. Translation invariance of the cube form gives the first inequality in (IV.6.5). The second follows from $L \leq B_k$ and $p/M = w \geq m/(3k)$. $\square$

Corollary IV.6.8 (Explicit strict nonuniformity coefficient). *Under Theorem IV.6.7, the coefficient*

$$c_k = \frac{1}{18k^2 2^k (2k)^k R_{k+1}(k2^{k-1})}$$

gives strict failure of the inequality $\sum_i \mathcal{U}_{k+1}(1_{R_i} f e_p(-P)) \leq c_k \alpha m^{k+2} M$. For $k \geq 2$,

$$\log_2(1/c_k) \leq k^2 + k \log_2 k + (2 + \log_2 e)k + 2 \log_2 k + \log_2(9e).$$

Proof. The lower bound in (IV.6.5) is twice the proposed threshold and is positive. The logarithmic estimate follows from $R_{k+1}(Q_k) \leq (eQ_k/(k+1))^{k+1} \leq e^{k+1} 2^{k^2-1}$. $\square$

This gives a local coefficient of size $2^{-k^2 - k \log_2 k - O(k)}$ at the stage where Proposition 17.7 uses $2^{-2(k+1)^3}$. The statement above specifies the field characteristic and the multi-affine extension used in the phase argument. It does not recompute the later inverse-theorem induction or the final bound for $r_k(N)$. Those are separate problems, and the local improvement alone does not supply a new global quantitative Szemerédi theorem.

[write] 6 October 2026. *Compared with the formal statement* (checked at HEAD): `proposition_17_7` (Sections17_18.lean, line 140; exact companion `proposition_17_7_holds`, `Proofs17LocalizedPhaseRemoval.lean`, line 2900) has the coefficient $2^{-2(k+1)^3}$. Its proof already bounds the profiles by a joint region count (`hyperplaneRegionCount`), as said at the start of this section. Source 01's Proposition IV.3.3 is a weaker selection interface at the same step: its loss is B_k^2 with source 01's B_k , against $2^k B_k$ here. Not formalized; nothing is propagated to Theorem 18.1, which is open in the ledger.

[write] 6 October 2026; sources 06–10. *A third coefficient (source 06)*. Source 06's Theorem IV.8.1 proves the localized phase removal with $c_k = 1/(8K_k^2)$, $K_k = (2(k+1))^k \mathcal{R}_{k+1}(k2^{k-1})$ (source 01's B_k), so $\log_2(1/c_k) = 2k^2 + O(k \log k)$ (Corollary IV.8.2). It keeps the full L^2 loss over pattern classes, as source 01's selection interface does, and is weaker than Corollary IV.6.8. At $k = 8$ (recomputed at the second write): $\log_2(1/c_8) \approx 121.66$ here, ≈ 212.70 in source 06, and $2(k+1)^3 = 1458$ in `proposition_17_7`. Not formalized.

[write] 6 October 2026; sources 15–20. *A stronger coefficient and a like-for-like comparison* (re-derived at the fourth write). Source 16's Theorem F.21.1 proves Proposition 17.7 with $c_k = (4/(3(k+2)))^{k+2}$, a loss of order $k \log k$ (Section IV.10). The comparison needs one normalization. Source 04's Corollary IV.6.8 normalizes by the input width, m^{k+2} , with cells of length q or $q+1$, $q = \lceil h/k \rceil$; the formal `proposition_17_7`, source 06's Theorem IV.8.1 and source 16 normalize by the largest cell length $l+1$. Since $m^{k+2} = (m/(q+1))^{k+2}(q+1)^{k+2}$, that corollary gives the coefficient $c_k(m/(q+1))^{k+2}$ in the $(l+1)$ normalization, with $l = q$. At $k = 8$ its loss is 91.07 bits uniformly over admissible m (the worst case is $m = 25$) and tends to $121.66 - 10 \log_2 16 = 81.66$ bits as $m \rightarrow \infty$; using the first inequality of Theorem IV.6.7 with $w \geq q$ in place of $w \geq m/(3k)$ gives 89.19 and 80.49. So the 121.66 bits printed above, and in the note at the end of Section IV.8, are in that corollary's own normalization and are not comparable with 212.70 (source 06) and 1,458 (formal), which use the $(l+1)$ one; the ordering stated there is unchanged. In one normalization, at $k = 2, 4, 8$: source 16 loses 6.34, 13.02, 29.07 bits; source 04's corollary 11.19, 29.37, 91.07; source 06 21.15, 65.35, 212.70; the formal statement 54, 250, 1,458 (the intake of batch 115 compared the same 121.7 bits; the intake of batch 117 found the difference of normalizations). Every value was recomputed at the write in exact arithmetic over every odd m with $3k < m < 20,000$; for larger m , $q+1 \leq (m+4k-1)/(2k)$ gives $m/(q+1) \geq 2km/(m+4k-1)$, which increases with m , so the loss stays below its bound at $m = 20,000$ (below 82 bits at $k = 8$).

IV.7 Exact floor-pattern counts and asymptotically sharp carry entropy

Source 06, delivered Section 2.

We first prove a slightly more general result than the Boolean-cube statement. Let $V \subset \mathbb{Z}_{\geq 0}^k$ be finite and contain 0 and the coordinate vectors $e_1, \dots, e_k$. Put

$$h(V) = \sum_{v \in V \setminus \{0\}} (\|v\|_1 - 1), \quad \tilde{h}(V) = \sum_{v \in V} \|v\|_1.$$

These are counts with multiplicity of the possible integer-level hyperplanes in an open unit cube; coincident hyperplanes can only improve the resulting upper bounds.

Definition IV.7.1 (Carry families). For $t \in [0, 1)^k$ define $c_t : V \rightarrow \mathbb{Z}$ by $c_t(v) = \lfloor v \cdot t \rfloor$. Let

$$C(V) = \#\{c_t : t \in [0, 1)^k\}, \quad C(V; M) = \#\{c_t \bmod M : t \in [0, 1)^k\}.$$

For $t_0 \in [0, 1)$ also define $d_{t_0, t}(v) = \lfloor t_0 + v \cdot t \rfloor$, and let $D(V)$ and $D(V; M)$ be the corresponding cardinalities. For $V = \{0, 1\}^k$ abbreviate these quantities as $C_k, C_{k, M}, D_k, D_{k, M}$.

The word “carry” describes precisely the integer part remaining after the integer parts of the coefficients have been removed. All these families are finite, since every coordinate takes values in a finite integer interval.

Lemma IV.7.2 (One-sided removal of boundary points). *Every carry pattern c_t is realized by a point of $(0, 1)^k$ for which $v \cdot t$ is nonintegral for all $v \in V \setminus \{0\}$. Every affine carry pattern is likewise realized in $(0, 1)^{k+1}$ away from all its integer-level hyperplanes.*

Proof. Replace t by $t + \eta(1, \dots, 1)$ for a sufficiently small $\eta > 0$. The new point remains below 1 in every coordinate. For every nonzero $v \in V$, the sum increases by $\eta \|v\|_1 > 0$. Choose η small enough that it is less than the distance to the next integer, divided by $\|v\|_1$, for every such v . A sum initially at an integer moves to the side having that same floor; every other sum remains between its original neighboring integers. The same argument for the coefficient vectors $(1, v)$ proves the affine assertion. Finiteness makes a common positive choice of η possible. $\square$

Theorem IV.7.3 (Exact separation of range and carries). *Let k, r, M be positive integers. The number of functions*

$$v \mapsto \lfloor v \cdot \alpha \rfloor, \quad \alpha \in (-r, r)^k,$$

is exactly $(2r)^k C(V)$. The number after reduction modulo M is exactly

$$\min(2r, M)^k C(V; M). \quad (\text{IV.7.1})$$

Moreover,

$$C(V; M) \leq C(V) \leq \mathcal{R}_k(h(V)), \quad D(V; M) \leq D(V) \leq \mathcal{R}_{k+1}(\tilde{h}(V)). \quad (\text{IV.7.2})$$

If $M \geq \max_{v \in V} \|v\|_1$, then $C(V; M) = C(V)$.

Proof. Write $\alpha = q + t$ with $q_i = \lfloor \alpha_i \rfloor$ and $0 \leq t_i < 1$. Then $q_i \in \{-r, \dots, r-1\}$ and

$$\lfloor v \cdot \alpha \rfloor = v \cdot q + c_t(v).$$

At $v = e_i$ the carry is zero, so the resulting pattern recovers q_i ; subtracting $v \cdot q$ then recovers c_t . Conversely every q in the stated integer box and every carry pattern are simultaneously realizable: use Lemma IV.7.2 to choose all $0 < t_i < 1$. In particular $q_i = -r$ causes no endpoint difficulty. This proves exact factorization.

Modulo M , the coordinate evaluations recover the residues $q_i \bmod M$. There are exactly $\min(2r, M)$ available residues in each coordinate. After these are recovered, the residual carry modulo M is recovered as well. Conversely all pairs are realizable, proving (IV.7.1). The possible nonzero carry values lie between 0 and $\|v\|_1 - 1$; hence reduction is injective on these coordinate ranges under the final hypothesis on M .

Inside the open unit cube, the only relevant hyperplanes for c_t are $v \cdot t = j$ with $1 \leq j \leq \|v\|_1 - 1$. Their number is at most $h(V)$. Each pattern is realized off these hyperplanes by Lemma IV.7.2, and each open region has a constant pattern. In fact a fixed pattern defines an intersection of strict linear inequalities, so it cannot split into disconnected pieces in the cube. The arrangement bound proves the first estimate. For affine carries the thresholds are $t_0 + v \cdot t = j$, $1 \leq j \leq \|v\|_1$, giving $\tilde{h}(V)$ hyperplanes in dimension $k + 1$. $\square$

The nonnegativity assumption on V has a specific purpose. If both v and $-v$ are allowed, a parameter with $v \cdot t = 0$ gives floors $(0, 0)$ that a generic perturbation need not preserve. Thus an extension to arbitrary signed coefficient families requires explicit treatment of lower-dimensional faces; the preceding proof does not silently cover that setting.

[write] 6 October 2026; sources 06–10. Theorem IV.7.3 for $V = \{0, 1\}^k$ is source 01’s Theorem IV.1.2 and source 04’s Theorem IV.6.3; the modular count (IV.7.1) is source 04’s modular count; the general family V corresponds to source 01’s Proposition IV.3.4, which bounds by the rank of the feature span instead of k (note there). Source 01’s example before its Section IV.4 is the signed-family obstruction described in the last paragraph. The bound $C_k \leq \mathcal{R}_k(h_k)$ is $C_k \leq \mathcal{R}(H_k, k) = R_k(H_k)$ of sources 01 and 04.

IV.7.1 The Boolean cube and the improvement of Corollary 17.5

For the Boolean cube, counting by support size gives

$$h_k := h(\{0, 1\}^k) = \sum_{s=2}^k \binom{k}{s} (s-1) = (k-2)2^{k-1} + 1, \quad (\text{IV.7.3})$$

$$\tilde{h}_k := \tilde{h}(\{0, 1\}^k) = \sum_{s=1}^k \binom{k}{s} s = k2^{k-1}. \quad (\text{IV.7.4})$$

The formula for h_k also gives $h_1 = 0$. Consequently,

$$\#\{\varepsilon \mapsto \lfloor \varepsilon \cdot \alpha \rfloor : \alpha \in (-r, r)^k\} = (2r)^k C_k \leq (2r)^k \mathcal{R}_k(h_k). \quad (\text{IV.7.5})$$

This replaces the source’s 2^{2rk^3} upper bound by a polynomial in r with the exact exponent and exact coefficient C_k . For example, $C_1 = 1$, $C_2 = 2$, and $C_3 = 10$. For $k = 3$, if the total fractional sum is below 1 all pair carries vanish, and if it is at least 2 all pair carries are 1. When the total is between 1 and 2, all eight patterns of the three pair carries occur. To realize respectively zero, one, two, or three positive pair carries in this middle slab, use $(.4, .4, .4)$, $(.6, .6, .1)$, $(.8, .3, .3)$, or $(.6, .6, .6)$ and permute coordinates. This proves the count $1 + 8 + 1 = 10$ directly.

The modular formula is exact even when the modulus is small:

$$\#\{\lfloor \varepsilon \cdot \alpha \rfloor \bmod M\} = \min(2r, M)^k C_{k,M}. \quad (\text{IV.7.6})$$

In particular, once $2r \geq M$ no further range loss is necessary. For $M \geq k$ the coefficient is the unreduced C_k .

IV.7.2 Affine patterns and translation classes

For $\alpha_0 \in \mathbb{R}$ and $\alpha \in (-r, r)^k$, consider

$$F_{\alpha_0, \alpha}(\varepsilon) = \lfloor \alpha_0 + \varepsilon \cdot \alpha \rfloor \pmod{M}.$$

Two such maps are equivalent if they differ by a constant in $\mathbb{Z}/M\mathbb{Z}$. Let $L_k(r, M)$ be the number of equivalence classes that occur.

Theorem IV.7.4 (Affine and modular carry bound). *For positive integers k, r, M ,*

$$L_k(r, M) \leq \min(2r, M)^k D_{k,M} \leq \min(2r, M)^k \mathcal{R}_{k+1}(k2^{k-1}). \quad (\text{IV.7.7})$$

The number of distinct maps is exactly $ML_k(r, M)$, and hence is at most $M \min(2r, M)^k \mathcal{R}_{k+1}(k2^{k-1})$.

Proof. Write $\alpha_0 = q_0 + t_0$ and $\alpha_i = q_i + t_i$ with fractional parts in $[0, 1)$. Then

$$F_{\alpha_0, \alpha}(\varepsilon) = q_0 + \varepsilon \cdot q + d_{t_0, t}(\varepsilon) \pmod{M}.$$

The constant q_0 disappears on passing to translation classes; the remaining choices give the upper bound. Unlike the nonaffine case, this decomposition need not be injective, because affine carries at singletons can be 1. Finally the family is invariant under addition of every constant modulo M , and that action is free: evaluation at $\varepsilon = 0$ detects the constant. Each orbit therefore has exactly M maps. $\square$

IV.7.3 Matching leading-order entropy

The next result determines the scale of the carry coefficient rather than merely improving an upper bound. The one external asymptotic input is the theorem of Zuev that the number $T(n)$ of Boolean linear threshold functions on $\{0, 1\}^n$ satisfies

$$\log_2 T(n) = (1 + o(1))n^2 \quad (n \rightarrow \infty) \quad [9, 8]. \quad (\text{IV.7.8})$$

A strict threshold representation can always be used: a separator on a finite cube can be moved slightly so that no value is zero.

Lemma IV.7.5 (Embedding thresholds into carries). *Let $T_+(n)$ denote the number of threshold functions admitting a representation with nonnegative coordinate weights. Then*

$$T(n) \leq 2^n T_+(n), \quad D_{n, M} \geq T_+(n) - 1, \quad C_{n+1, M} \geq D_{n, M} \quad (M \geq 2).$$

The last two inequalities also hold without modular reduction.

Proof. Flip each input coordinate whose coefficient in a chosen threshold representation is negative. The resulting function has nonnegative weights. Recording the subset of flipped coordinates and the resulting function recovers the original function. There are at most 2^n choices of the subset, proving the first inequality.

A nonconstant such function can be written $g(\varepsilon) = \mathbf{1}_{\{w \cdot \varepsilon > \theta\}}$ with all $w_i > 0$ and $0 < \theta < \sum_i w_i$. Zero weights can be perturbed to small positive ones because the finite strict separator has a positive margin. Choose $\lambda > 0$ sufficiently small that

$$t_0 = 1 - \lambda\theta \in (0, 1), \quad t_i = \lambda w_i \in (0, 1), \quad \sum_i t_i < 1.$$

All affine sums lie in $(0, 2)$, and

$$\lfloor t_0 + \varepsilon \cdot t \rfloor = g(\varepsilon).$$

The constant-zero function is also realizable by small positive parameters. Only the constant-one function must be omitted, because an affine carry at the zero vertex is 0. Distinct Boolean outputs remain distinct modulo every $M \geq 2$. This proves $D_{n, M} \geq T_+(n) - 1$. Finally restricting a nonaffine carry in $n + 1$ variables to the face whose first coordinate is 1 produces an arbitrary affine carry in n variables. This restriction is surjective and gives the final bound. $\square$

Theorem IV.7.6 (Asymptotically sharp carry entropy). *As $k \rightarrow \infty$,*

$$\log_2 C_k = (1 + o(1))k^2, \quad \log_2 D_k = (1 + o(1))k^2. \quad (\text{IV.7.9})$$

Both formulas remain valid for $C_{k,M}$ and $D_{k,M}$, uniformly for integer choices $M = M(k) \geq 2$. In particular, for fixed positive integer r the exact nonaffine count is

$$(2r)^k 2^{(1+o(1))k^2}.$$

For fixed $r \geq 1$ and $M \geq 2$, $\log_2 L_k(r, M) = (1 + o(1))k^2$ as well.

Proof. Equations (IV.7.2)–(IV.7.4) and (F.11.2) give $\log_2 C_k \leq k^2 + O(k)$ and $\log_2 D_k \leq k^2 + O(k)$; the finitely many small dimensions where $m < d$ can be treated directly. Lemma IV.7.5 and (IV.7.8) give

$$C_{k,M} \geq 2^{-(k-1)}T(k-1) - 1, \quad D_{k,M} \geq 2^{-k}T(k) - 1.$$

These lower bounds are independent of $M \geq 2$, while the upper bounds are the unreduced ones, proving the uniform assertion. For translation classes, the threshold patterns from the lemma have value 0 at the zero vertex. Two distinct such patterns cannot differ by a constant modulo $M \geq 2$. They are realized with all slopes in $(0, 1) \subset (-r, r)$, so $L_k(r, M) \geq T_+(k) - 1$. Combine this with Theorem IV.7.4. $\square$

Thus the cubic dimension exponent of the coarse source count can be replaced by a quadratic one, but a uniform subquadratic exponent for these pattern families is impossible. This lower bound concerns the number of patterns, not an optimality assertion for every proof of localized phase removal. A proof that avoids enumerating every pattern could conceivably have smaller analytical losses.

[write] **6 October 2026; sources 06–10.** Theorem IV.7.6 is source 01’s Theorem IV.2.3 (same external input, Zuev’s count), extended uniformly in the modulus and to translation classes; source 01 also bounds the secondary term from below by $-k \log_2 k - O(k)$. Together they answer source 04’s Question 8 (note there). Source 04’s elementary lower bound $2^{k(k-1)/2}$ (Theorem IV.6.2) is a third, non-asymptotic lower bound.

IV.8 A quadratic-exponent loss in localized phase removal

Source 06, delivered Section 4.

The carry bound has a direct analytical consequence. We give the argument because a better combinatorial count alone does not justify changing a downstream constant. The following formulation makes the partition scale explicit and supplies a replacement for the local loss in the source’s Proposition 17.7.

Let p be prime and write $e_p(t) = \exp(2\pi it/p)$. For a function on $\mathbb{F}_p$ define its unnormalized cube energy by

$$\mathcal{E}_d(g) = \sum_{s \in \mathbb{F}_p} \sum_{a \in \mathbb{F}_p^d} \prod_{\varepsilon \in \{0,1\}^d} \mathcal{C}^{|\varepsilon|} g(s - \varepsilon \cdot a) = p^{d+1} \|g\|_{U^d}^{2^d}.$$

Here $\mathcal{C}$ denotes complex conjugation. This is a nonnegative real number. For a partition $\mathcal{Q} = (Q_i)_{i=0}^{M-1}$ put

$$\mathcal{E}_d(g; \mathcal{Q}) = \sum_i \mathcal{E}_d(\mathbf{1}_{Q_i} g).$$

If every cell has length at most ℓ , the partition-uniformity normalization at degree $d-1$ is $M\ell^{d+1}$. Define the explicit integer

$$K_k = [2(k+1)]^k \mathcal{R}_{k+1}(k2^{k-1}), \quad c_k = \frac{1}{8K_k^2}. \quad (\text{IV.8.1})$$

The letter K_k here denotes a pattern bound; it is unrelated to the large exponents in the polynomial-partition section of the source.

Theorem IV.8.1 (Localized phase removal with the carry bound). *Let $k \geq 1$, let $p > k + 1$ be prime, and let m be odd with $1 \leq m \leq \sqrt{p}$. Let*

$$P = \prod_{i=1}^k (r_i + I), \quad I = \{-h, \dots, h\}, \quad 2h + 1 = m.$$

All sets are viewed in $\mathbb{F}_p$. Suppose $|f| \leq 1$, $\alpha > 0$, and σ is the restriction to P of a polynomial separately of degree at most one in each of its k variables. Assume

$$\sum_{x \in P} \left| \sum_{s \in \mathbb{F}_p} \prod_{\varepsilon \in \{0,1\}^k} c^{|\varepsilon|} f(s - \varepsilon \cdot x) e_p(-s\sigma(x)) \right|^2 \geq \alpha p^2 m^k. \quad (\text{IV.8.2})$$

Then there is a polynomial ϕ of degree at most $k + 1$ and a partition $\mathcal{Q}$ into modular intervals, each of length at least $m/(3k)$ and at most some integer ℓ , such that

$$\mathcal{E}_{k+1}(f e_p(-\phi); \mathcal{Q}) > c_k \alpha M \ell^{k+2}. \quad (\text{IV.8.3})$$

The same assertion holds for input progressions with any common nonzero step, with modular progressions in the conclusion.

Proof. Small boxes. If $m \leq 3k$, use the singleton partition and $\phi = 0$. For each x , Cauchy–Schwarz in s bounds the Fourier square in (IV.8.2) by

$$p \sum_s \prod_{\varepsilon \in \{0,1\}^k} |f(s - \varepsilon \cdot x)|^2.$$

Arithmetic–geometric mean, followed by translation invariance of the s -sum, bounds this by $p \sum_s |f(s)|^{2^{k+1}}$. Thus the hypothesis implies $\sum_s |f(s)|^{2^{k+1}} \geq \alpha p$. This is exactly the singleton-partition cube energy, while $M = p$, $\ell = 1$, and $c_k < 1$. The singleton length is at least $m/(3k)$, as required.

The partition and its pattern classes. Suppose $m > 3k$. Put

$$q = \lceil h/k \rceil, \quad M = \lfloor p/q \rfloor, \quad w = p/M, \quad \ell = q + 1,$$

and, using integer representatives, set

$$Q_j = \{s \in \mathbb{F}_p : jw \leq s < (j+1)w\}, \quad 0 \leq j < M.$$

Since $q \leq \sqrt{p}$, one has $q \leq M$. Writing $p = qM + b$ gives $0 \leq b < q \leq M$, so the cell lengths are q or $q + 1$. Also

$$q \geq h/k \geq m/(3k), \quad w \geq q \geq \ell/2, \quad m \geq \ell. \quad (\text{IV.8.4})$$

Set $f_\varepsilon(s) = f(s - \varepsilon \cdot r)$ and $\tau(a) = \sigma(a + r)$, taking the indicated global multiaffine extension. In the expansion obtained by partitioning each f_ε , a nonzero contribution at (s, a) has cell labels

$$j(\varepsilon) = \lfloor s/w - \varepsilon \cdot (a/w) \rfloor \pmod{M}.$$

This identity is well defined for representatives because $p = Mw$. Moreover $|a_i/w| \leq h/q \leq k < k+1$. By Theorem IV.7.4, the number L of possible label functions up to simultaneous translation is at most K_k .

Keeping the pattern loss explicit. Choose one representative j_ν of each translation class and define

$$T_{\nu,i}(a) = \sum_s e_p(-s\tau(a)) \prod_{\varepsilon \in \{0,1\}^k} \mathcal{C}^{|\varepsilon|}(\mathbf{1}_{Q_{j_\nu(\varepsilon)+i}} f_\varepsilon)(s - \varepsilon \cdot a).$$

Indices on cells are modulo M . The original inner sum is $\sum_\nu \sum_i T_{\nu,i}(a)$. Cauchy–Schwarz over the L classes, selection of one class, and then Cauchy–Schwarz over its M translates give a ν such that

$$\sum_{i=0}^{M-1} \sum_{a \in I^k} |T_{\nu,i}(a)|^2 \geq \frac{\alpha p^2 m^k}{L^2 M}. \quad (\text{IV.8.5})$$

Extend the a -sum to $\mathbb{F}_p^k$; every added term is nonnegative.

Removing the multiaffine phase. Expanding each square and putting $y = s - t$ makes its phase $e_p(-y\tau(a))$. Every monomial in $y\tau(a)$ is square-free and has degree at most $k+1$. For a square-free monomial on a nonempty coordinate set J , finite differences give

$$|J|! \prod_{j \in J} z_j = \sum_{\varepsilon \in \{0,1\}^J} (-1)^{|\varepsilon|} (s - \varepsilon \cdot z)^{|J|}.$$

All the factorials are invertible because $p > k+1$. Regard this sum as a sum on the full cube by assigning zero to unused vertices, and add the identities for the monomials. We obtain polynomials ϕ_η of degree at most $k+1$ such that

$$y\tau(a) = \sum_{\eta \in \{0,1\}^{k+1}} (-1)^{|\eta|} \phi_\eta(s - \eta \cdot (a, y)).$$

Consequently each expanded square in (IV.8.5) is a mixed cube form in the functions

$$\mathbf{1}_{Q_{j_\nu(\eta')+i}} f_{\eta'} e_p(-\phi_\eta), \quad \eta' = (\eta_1, \dots, \eta_k).$$

Apply the Gowers–Cauchy–Schwarz inequality to this form and then arithmetic–geometric mean to the resulting 2^{k+1} energy factors. Summing in i and averaging in η shows that for some η , with $g = f_{\eta'} e_p(-\phi_\eta)$,

$$\sum_i \mathcal{E}_{k+1}(\mathbf{1}_{Q_i} g) \geq \frac{\alpha p^2 m^k}{L^2 M}. \quad (\text{IV.8.6})$$

The shift of the cell index by $j_\nu(\eta')$ disappears because it permutes the M cells. The Gowers–Cauchy–Schwarz inequality used here is the standard cube inequality proved in Section 3 of the source; it applies to all complex functions, including these restrictions.

Divide (IV.8.6) by $M\ell^{k+2}$. Since $p/M = w$, (IV.8.4) yields

$$\frac{\mathcal{E}_{k+1}(g; \mathcal{Q})}{M\ell^{k+2}} \geq \frac{\alpha}{L^2} \left(\frac{w}{\ell}\right)^2 \left(\frac{m}{\ell}\right)^k \geq \frac{\alpha}{4K_k^2} > c_k \alpha.$$

Finally translate the partition by $-\eta' \cdot r$ and replace $\phi_\eta(s)$ by $\phi_\eta(s + \eta' \cdot r)$ to express g as a polynomial twist of the original f . A translation does not change cell lengths or cube energies. For a common nonzero input step, first dilate all variables to make that step 1; multiplication by the step is a bijection on $\mathbb{F}_p$, and the Fourier phase and the polynomial change by the same invertible substitution. $\square$

[write] 6 October 2026; sources 15–20. Source 16’s Theorem F.21.1 has the same normalization (cells of against 212.70 here (note after Corollary IV.6.8).

Corollary IV.8.2 (Size of the new local loss). *The constant in (IV.8.3) satisfies*

$$\log_2(c_k^{-1}) = 3 + 2k \log_2(2(k+1)) + 2 \log_2 \mathcal{R}_{k+1}(k2^{k-1}) = 2k^2 + O(k \log k).$$

Thus the explicit local loss has a quadratic leading exponent, compared with the cubic exponent $2(k+1)^3$ in the displayed source formulation.

The last equality follows directly from the binomial-sum upper bound and its largest term; alternatively, upper and lower bounds of order $k^2 + O(k \log k)$ for its logarithm suffice. The sharper raw estimate (IV.8.6) can be retained when the actual number of classes or the observed cell-size ratios are known. No asymptotic approximation is needed to use the explicit constant (IV.8.1).

This proof accounts for the scale of every partition cell and for the factor L^2 . Later structural extraction and stopping thresholds are separate inputs. Improving this local loss does not, on its own, justify substituting a new exponent into the final Szemerédi theorem.

[write] 6 October 2026; sources 06–10. *Compared with the other two coefficients at this step.* Source 04’s Theorem IV.6.7 with Corollary IV.6.8 has the stronger coefficient $2^{-k^2 - k \log_2 k - O(k)}$ (fixed-slope profiles, loss $2^k B_k$ in place of K_k^2); source 01’s Proposition IV.3.3 is a selection interface with the same L^2 loss as here. At $k = 8$: about $2^{-121.66}$ (04), $2^{-212.70}$ (here; recomputed at the second write), and 2^{-1458} in the formal `proposition_17_7` (exact companion `proposition_17_7_holds`). The numerical lemma “ $L^2 c < 1/4$ ” cited in Section V.26 is in `Proofs17LocalizedPhaseRemoval.lean`. Not formalized.

[write] 6 October 2026; sources 15–20. The $2^{-121.66}$ above is in source 04’s own normalization (by the input width m^{k+2}); in the $(l+1)$ normalization of the other two values source 04’s coefficient loses 91.07 bits at $k = 8$, and source 16’s 29.07 (note after Corollary IV.6.8).

IV.9 Exact computations and reproducible verification

Source 06, delivered Section 7.

The general theorems do not depend on experimental extrapolation. Computations serve two narrower purposes: certifying small carry counts that are useful examples, and checking the quantitative inequalities on finite configurations that exercise their normalizations and endpoints.

IV.9.1 Certified carry counts

The included certificate package gives the following values. The zero-dimensional entries use the natural empty-cube convention.

k	C_k	$\mathcal{R}_k(h_k)$	D_k	$\mathcal{R}_{k+1}(\tilde{h}_k)$
0	1	1	1	1
1	1	1	2	2
2	2	2	6	15
3	10	26	38	794
4	154	3214	not computed	not tabulated

The entry “not computed” describes this package’s scope; it is not a claim that the value is unknown in the literature. A short initial sequence match is not sufficient to identify these objects with an existing integer sequence.

[write] **6 October 2026; sources 06–10.** $C_1, \dots, C_4$ agree with source 01’s certified values (Proposition IV.5.1), so $C_4 = 154$ is certified twice with different formats; D_1, D_2, D_3 agree with source 04’s certified A_1, A_2, A_3 . The intake reran the verifier (all 215 witnesses and 1,165 dual certificates exact). C_5 and D_4 are not certified by any source; the intake’s floating-point values $C_5 = 8410$ and $A_4 = 770$ are item IV.2 of Section V.2.

The enumeration adds the relevant hyperplanes one at a time and branches on their two strict sides. For a sign prefix in dimension d , encode the unit-cube and selected side constraints as

$$A_i \cdot x + t \leq b_i.$$

A rational point x with a rational $t > 0$ proves the existence of a strictly feasible chamber. To certify that a branch has no strict solution, the package supplies rational multipliers $y_i \geq 0$ satisfying

$$\sum_i y_i = 1, \quad \sum_i y_i A_i = 0, \quad \sum_i y_i b_i \leq 0.$$

Multiplying and summing the inequalities would force $t \leq \sum_i y_i b_i \leq 0$, a contradiction. This is an exact certificate, not a numerical tolerance test. The verifier checks the entire binary prefix tree: every branch either ends in such a certificate or reaches a strictly witnessed full pattern. It also checks that the resulting carry maps are distinct.

The final data contain 215 rational chamber witnesses and 1,165 exact nonnegative dual certificates across all tabulated cases. A numerical linear-programming solver is used only to propose rational certificates. The separate verifier uses Python’s standard library and `fractions.Fraction`; it verifies every relation exactly and rejects execution with Python’s assertion-disabling optimization flag. Lemma IV.7.2 is what makes this open-chamber enumeration complete for the original floor patterns, including their boundary values.

IV.9.2 Uniformity and density checks

The uniformity checker evaluates the finite-group progression counts and cube energies with integers and rational arithmetic. It raises the inequalities to integer powers instead of numerically approximating Gowers-norm roots. It checks all subsets of $\mathbb{Z}/5\mathbb{Z}$ for lengths 3, 4, and 5, all subsets of $\mathbb{Z}/7\mathbb{Z}$ for lengths 3 and 4, and eight specified subsets of $\mathbb{Z}/25\mathbb{Z}$ for lengths 3 and 4. All 368 set-length cases satisfy both the linear and quadratic bounds. It also checks 160 two-term cases and 168 three-term variance bounds, including 32 equality cases. The convenient sufficient existence criterion fires in 62 checked cases; every one contains a nonconstant progression.

The weighted extraction checker exercises 22,752 rational indicator/partition configurations under each of the ranges $[-1, 1]$ and $[-\delta, 1 - \delta]$, for 45,504 exact range cases, and verifies the two extremal constructions. These finite checks are supplementary to the proofs: unlike a carry-count certificate, a list of finite examples does not prove a universal inequality. The accompanying JSON files state the exact tested families and counts.

IV.9.3 Magnitude of the explicit improvements

Figure IV.9.1 evaluates the formulas proved above. In its left panel the new count bound is $2^k \mathcal{R}_k(h_k)$, so it uses the general arrangement estimate rather than small exact counts. The right panel uses

the fully explicit c_k from (IV.8.1).




Figure IV.9.1: Comparison of logarithmic bounds. Lower values mean a smaller count upper bound (left) or a larger guaranteed local correlation constant (right). These are comparisons of local statements, not plots of final Szemerédi thresholds.

For example, at $k = 4, r = 1$, the displayed floor-count exponent 128 is replaced by the arrangement-based exponent approximately 15.650, and the exact count is still smaller, namely $2^4 C_4 = 2464$. At $k = 8$, the displayed local phase loss is 2^{-1458} , whereas (IV.8.1) gives a constant approximately $2^{-212.704}$. These comparisons quantify genuine local savings while leaving the later structural dependencies explicit.

IV.10 Source 16's chapter: Proposition 17.7 without carry patterns [write]

Added in the fourth write, 6 October 2026.

The next four sections are source 16's Sections 2–5, on Gowers's Proposition 17.7, the localized phase removal that the floor-pattern counts of this Part serve. Source 16 localizes every cube to one sliding interval window and counts exactly how many windows contain it, so no carry pattern is counted. *Logically this chapter should open the Part's treatment of Proposition 17.7:* of the four coefficients printed here it is the strongest at every k , of order $2^{-k \log_2 k - O(k)}$ against $2^{-k^2 + O(k)}$ for source 04 in the same normalization, $2^{-2k^2 - O(k \log k)}$ for source 06 and $2^{-2(k+1)^3}$ in the paper and the formal statement. It is printed here, at the end of Part IV, because no number printed earlier may move. Its introduction, with the main theorem (Theorem F.21.1), is Section F.21; its other sections are in Parts I, III and V. Its notation is in the table at the head of this Part.

What is proved, in one paragraph. Under exactly the hypotheses of Proposition 17.7, a polynomial phase of degree at most $k + 1$ and a partition into proper modular progressions of lengths l or $l + 1$, $m/(3k) \leq l$ and $l + 1 \leq 3km$, on which uniformity fails at $\alpha(4/(3(k+2)))^{k+2}$, normalized by $(l+1)^{k+2}$ (Theorem F.21.1, from the exact window estimate, Theorem IV.13.2); the width optimization gives $(2/(k+2))^{k+2}$ for fixed k (Proposition IV.13.4), and no conclusion with cells of size proportional to m can beat $2^{k+1}/(k+2)!$, so the order $k \log k$ is optimal in that class (Proposition IV.13.5). The polarization is source 04's and the formal development's; the factorial criterion and the $\mathbb{F}_p$ functional test are new in this report (Theorems IV.12.2, IV.12.4). Section IV.14 keeps the original small cells

at the price of a global count; it is source 04’s Lemmas IV.6.5–IV.6.6.

The four coefficients in one normalization (re-derived at the write; bits lost, $-\log_2$ of the coefficient, with cells of length l or $l + 1$ normalized by $(l + 1)^{k+2}$, as in `proposition_17_7`).

k	Source 16	Source 04, converted	Source 06	Formal statement
2	6.34	11.19	21.15	54
4	13.02	29.37	65.35	250
8	29.07	91.07	212.70	1,458

Source 04’s Corollary IV.6.8 normalizes by the input width m^{k+2} ; converted to $(q + 1)^{k+2}$, $q = \lceil h/k \rceil$, it loses the values shown uniformly over admissible m , and 81.66 bits as $m \rightarrow \infty$ at $k = 8$; its own normalization gives the 121.66 bits quoted in the notes of the second write (note after that corollary). Source 01’s selection interface is weaker than source 04’s. *Caveat*: under the printed reading “each of length at most m ” (source 04’s), source 16’s cells, of length about $(k + 2)m/2$, are too long; source 16 follows the formal $(l + 1)$ reading, which Section 18 needs (it uses only $l \geq m/(3k)$ and lengths l or $l + 1$). Larger cells must be transported in Section 18; source 16’s remark after the proof of its theorem gives a sufficient condition.

Relation to the formal statements (checked at HEAD 0b4890a1c). `proposition_17_7` (`Section17_18.lean`, line 140) is an exact companion (`proposition_17_7_holds`, `Proofs17LocalizedPhaseRemoval.lean`, line 2900), with the coefficient $2^{-2(k+1)^3}$; source 16’s theorem implies it (note in Section F.21), and is not formalized. The carry counts stay needed for `corollary_17_5` and `corollary_17_6` (exact companions, `Proofs17FinitePatterns.lean`, lines 396, 501). No formal declaration consumes `proposition_17_7`; `theorem_18_1` is open.

IV.11 Cube forms, normalizations, and finite partitions

Source 16, delivered Section 2.

IV.11.1 Normalized and unnormalized forms

On a finite abelian group G , all expectations are with respect to uniform probability. If $u : G \rightarrow \mathbb{C}$, set

$$\|u\|_{U^j}^{2^j} = \mathbb{E}_{x, h_1, \dots, h_j} \prod_{\epsilon \in \{0,1\}^j} \mathbb{C}^{|\epsilon|} u(x - \epsilon \cdot h). \quad (2.1)$$

For $j \geq 1$, this expression is nonnegative. For $j \geq 2$ its 2^j -th root is a norm. Its unnormalized version is

$$\mathcal{C}_j(u) = \sum_{h_1, \dots, h_{j-1}} \left| \sum_x \Delta(u; h_1, \dots, h_{j-1})(x) \right|^2. \quad (2.2)$$

Expanding the square and setting the last increment equal to the difference of the two base points proves $\mathcal{C}_j(u) = |G|^{j+1} \|u\|_{U^j}^{2^j}$. Formula (2.2) will be essential: restriction of the increment domain can safely be removed *before* expansion, because each term is a nonnegative square.

Lemma IV.11.1 (Mixed cube inequality). *For functions $u_\eta : G \rightarrow \mathbb{C}$, indexed by $\eta \in \{0,1\}^j$,*

$$\left| \sum_{x, h} \prod_{\eta} \mathbb{C}^{|\eta|} u_\eta(x - \eta \cdot h) \right| \leq \prod_{\eta} \mathcal{C}_j(u_\eta)^{1/2^j} \leq 2^{-j} \sum_{\eta} \mathcal{C}_j(u_\eta). \quad (2.3)$$

Proof. The first inequality is the Gowers–Cauchy–Schwarz inequality, with the normalization in (2.1). One obtains it by applying Cauchy–Schwarz successively in the j direction variables, at each step separating the two faces in that direction and duplicating the variables on the opposite face. After all directions have been processed, the factor associated with a vertex is its diagonal j -cube form. Each factor has exponent 2^{-j} . The total normalizing factor is $|G|^{j+1}$, the same on both sides. This is the finite cube inequality used in [1, Lemma 3.8]. The second inequality follows from the arithmetic–geometric mean inequality for 2^j nonnegative numbers. $\square$

A modular arithmetic progression is *proper* if its displayed parameterization is injective. A consecutive cyclic interval is allowed to cross the chosen representative $N - 1$; properness is a statement about distinct residues. When ordinary integer progressions are needed, a separate rectification statement is necessary.

For a partition $\mathcal{Q}$ with J cells and an integer U bounding all its cell sizes, define

$$\mathcal{B}_k(g; \mathcal{Q}, U) = \frac{\sum_{Q \in \mathcal{Q}} \mathcal{C}_{k+1}(\mathbf{1}_Q g)}{JU^{k+2}}. \quad (2.4)$$

The upper bound U is part of the notation. Changing U changes the quotient, even when the partition itself is unchanged.

IV.12 Explicit polarization and its algebraic limits

Source 16, delivered Section 3.

IV.12.1 The phases can be chosen before localization

Lemma IV.12.1 (Explicit multi-affine polarization). *Let R be a commutative ring in which $1!, \dots, (k+1)!$ are units. Write*

$$\tau(a) = \sum_{A \subseteq [k]} c_A \prod_{j \in A} a_j.$$

For $\epsilon \in \{0, 1\}^k$, put

$$\phi_\epsilon(X) = \sum_{A \supseteq \text{supp}(\epsilon)} \frac{c_A}{(|A| + 1)!} X^{|A|+1}. \quad (3.1)$$

Then, as a formal polynomial identity,

$$y\tau(a) = \sum_{\epsilon \in \{0,1\}^k} \sum_{b=0}^1 (-1)^{|\epsilon|+b} \phi_\epsilon(s - \epsilon \cdot a - by). \quad (3.2)$$

Thus the two phases at $(\epsilon, 0)$ and $(\epsilon, 1)$ coincide, and only 2^k candidate phases are needed.

Proof. For indeterminates $z_1, \dots, z_r$, successive backward differences give

$$\sum_{\omega \in \{0,1\}^r} (-1)^{|\omega|} \left(s - \sum_j \omega_j z_j \right)^r = r! \prod_j z_j. \quad (3.3)$$

One proof is to expand each power: a monomial missing some z_j cancels upon summing the two values of ω_j ; the only surviving monomial has every z_j to the first power, with coefficient $r!$. Apply

this identity with $r = |A| + 1$, with the variables a_j for $j \in A$ and y , multiply by $c_A/r!$, and sum over A . The allowed vertices in the term for A are exactly those ϵ supported inside A , with either value of b . This is (3.2). $\square$

The independence of these phases from all spatial cell labels is what allows localization followed by a single averaging argument. It also allows one to keep a sum over carry classes instead of selecting a class prematurely.

IV.12.2 An exact factorial criterion

The factorial hypothesis can be sharpened when individual coefficients have divisibility. The following statement is about *formal polynomial identities*; identities of functions on a finite ring require a separate analysis.

Theorem IV.12.2 (Coefficient criterion for vertex polarization). *Let R be a commutative ring and*

$$H(z_1, \dots, z_r) = \sum_{A \subseteq [r]} c_A \prod_{i \in A} z_i.$$

There exist one-variable polynomials $\psi_\eta \in R[X]$ satisfying

$$H(z) = \sum_{\eta \in \{0,1\}^r} (-1)^{|\eta|} \psi_\eta(s - \eta \cdot z) \quad (3.4)$$

as a formal identity if and only if

$$c_A \in |A|! R \quad \text{for every } A \subseteq [r]. \quad (3.5)$$

When this holds, the ψ_η may be chosen with degree at most r . For $H(a, y) = y\tau(a)$, the condition is $c_A \in (|A| + 1)! R$.

Proof. In the right side of (3.4), the coefficient of the square-free monomial $\prod_{i \in A} z_i$, with no s , receives contributions only from the degree- $|A|$ terms of the polynomials ψ_η . Each such contribution is a multiple of $|A|!$; its other factors are signs and zero-one coordinates of η . This proves necessity.

Conversely choose $b_A \in R$ with $c_A = |A|!b_A$. For each A , use (3.3) on just its coordinates and multiply by b_A . The case $A = \emptyset$ is the constant identity. Add all these identities and collect terms at each vertex. All degrees are at most r , proving sufficiency. $\square$

Corollary IV.12.3 (Cyclic coefficient test). *Over $\mathbb{Z}/N\mathbb{Z}$, condition (3.5) is equivalent to*

$$\gcd(N, |A|!) \text{ dividing an integer representative of } c_A.$$

In particular, $\gcd(N, r!) = 1$ suffices for every multi-affine H .

Proof. The image of multiplication by the integer a on $\mathbb{Z}/N\mathbb{Z}$ is precisely the set of residue classes divisible by $\gcd(a, N)$. Apply this with $a = |A|!$. $\square$

IV.12.3 A precise small-characteristic obstruction

Theorem IV.12.4 (Functional polarization over a prime field). *Let p be prime and let $H : \mathbb{F}_p^r \rightarrow \mathbb{F}_p$ be represented by the multi-affine polynomial above. It can be expressed in the form (3.4) as an identity of functions, allowing arbitrary functions $\psi_\eta : \mathbb{F}_p \rightarrow \mathbb{F}_p$, if and only if*

$$c_A = 0 \quad (|A| \geq p). \quad (3.6)$$

Proof. Every function $\mathbb{F}_p \rightarrow \mathbb{F}_p$ is represented by a polynomial of degree at most $p - 1$, by Lagrange interpolation. Substituting $s - \eta \cdot z$ produces a polynomial of total degree at most $p - 1$. Its degree in each variable is also at most $p - 1$.

A polynomial whose degree in each variable is less than p and which vanishes on the entire affine space has every coefficient zero: apply the univariate root bound in the last variable, and induct on the number of variables. Thus the reduced polynomial representing a function is unique. Since the multi-affine polynomial H is already reduced, it cannot contain a nonzero monomial of total degree at least p if it equals the indicated sum. This proves necessity. For sufficiency, all remaining factorials are nonzero in $\mathbb{F}_p$; apply (3.3) to every nonzero monomial. $\square$

For example, $z_1 \cdots z_p$ has no such representation, even with arbitrary one-variable $\mathbb{F}_p$ -valued functions at the vertices. This obstruction concerns classical field-valued phases. It makes no claim about phases with values in larger torsion subgroups of $\mathbb{R}/\mathbb{Z}$, or about other inverse-theorem formulations.

IV.13 Finite sliding-window localization without carry patterns

Source 16, delivered Section 4.

The proof in this section localizes all cube vertices to one sliding interval. A finite counting identity records exactly how many windows contain a cube. It replaces the enumeration of carry patterns in [1, Proposition 17.7]; the polarization phases remain those already constructed in Lemma IV.12.1.

IV.13.1 A finite family of sliding partitions

Let N be prime, let $h \geq 1$, put $R = kh$, and choose a positive integer M satisfying

$$R < w := \frac{N}{M} < N - R. \quad (4.1)$$

In particular $1 < M < N$. Set $U = \lfloor w \rfloor + 1 = \lceil w \rceil$. Embed $\mathbb{Z}/N\mathbb{Z}$ into $\mathbb{Z}/(NM)\mathbb{Z}$ by $x \mapsto Mx$. For $t \in \mathbb{Z}/(NM)\mathbb{Z}$, define

$$S_t = \{x \in \mathbb{Z}/N\mathbb{Z} : ((Mx - t) \bmod NM) \in \{1, \dots, N\}\}. \quad (4.2)$$

Thus one takes a block of N consecutive residues in the larger cyclic group and intersects it with the embedded copy of $\mathbb{Z}/N\mathbb{Z}$.

Lemma IV.13.1 (Sliding-window geometry). *Every S_t is a cyclic arithmetic progression with common difference one and cardinality $\lfloor w \rfloor$ or $\lceil w \rceil$. For every $j \in \{0, \dots, N - 1\}$, the family*

$$\mathcal{P}_j = (S_{j+\ell N})_{0 \leq \ell < M} \quad (4.3)$$

is a partition of $\mathbb{Z}/N\mathbb{Z}$. If $a \in [-h, h]^k \cap \mathbb{Z}^k$ and $s \in \mathbb{Z}/N\mathbb{Z}$, precisely

$$N - MD(a), \quad D(a) := \sum_{i=1}^k |a_i|, \quad (4.4)$$

of the sets S_t contain every vertex $s - \epsilon \cdot a$, $\epsilon \in \{0, 1\}^k$.

Proof. A block of N consecutive integers contains either $\lfloor N/M \rfloor$ or $\lceil N/M \rceil$ multiples of M . Their pullbacks are consecutive residues in $\mathbb{Z}/N\mathbb{Z}$. The M blocks starting at $j + \ell N + 1$ partition $\mathbb{Z}/(NM)\mathbb{Z}$, proving the partition assertion after intersection with the embedded subgroup.

The integer lifts of the cube vertices have span exactly $D(a)$: the maximum and minimum subset sums differ by $\sum_i |a_i|$. After embedding, their span is $MD(a) < N$. Write their extreme lifts as $u \leq v$. An integer block $\{t + 1, \dots, t + N\}$ contains the vertices exactly when $v - N \leq t \leq u - 1$, giving $N - (v - u) = N - MD(a)$ starts. There are no additional containing blocks obtained by traversing the opposite part of the circle. Indeed the complementary circular gap has length $NM - MD(a) > N$, by $w + R < N$; all gaps inside the displayed span are shorter than N . This also proves the assertion when some vertices coincide, including $D(a) = 0$. $\square$

IV.13.2 Exact phase extraction

Let $f_\epsilon : \mathbb{Z}/N\mathbb{Z} \rightarrow \mathbb{C}$, $\epsilon \in \{0, 1\}^k$, and let τ be multi-affine over $\mathbb{F}_N$, with $N > k + 1$. Define

$$F(a) = \sum_{s \in \mathbb{Z}/N\mathbb{Z}} e_N(-s\tau(a)) \prod_{\epsilon} \mathbb{C}^{|\epsilon|} f_\epsilon(s - \epsilon \cdot a), \quad E = \sum_{a \in [-h, h]^k} |F(a)|^2. \quad (4.5)$$

Integer increments in this formula are reduced modulo N . Let ϕ_ϵ be the polynomials in Lemma IV.12.1, and put $g_\epsilon = f_\epsilon e_N(-\phi_\epsilon)$.

Theorem IV.13.2 (Exact finite sliding-window estimate). *Under (4.1), some $\epsilon \in \{0, 1\}^k$ and $j \in \{0, \dots, N - 1\}$ satisfy*

$$\mathcal{B}_k(g_\epsilon; \mathcal{P}_j, U) \geq \frac{E(w - kh)^2}{N^2 U^{k+2}}. \quad (4.6)$$

In particular, if $m = 2h + 1$ and $E \geq \alpha N^2 m^k$, then

$$\mathcal{B}_k(g_\epsilon; \mathcal{P}_j, U) \geq \alpha \frac{m^k (w - kh)^2}{U^{k+2}}. \quad (4.7)$$

The phases are determined by τ before any window is selected.

Proof. Let $F_t(a)$ be the expression defining $F(a)$, with each f_ϵ replaced by $\mathbf{1}_{S_t} f_\epsilon$. Lemma IV.13.1 gives the exact identity

$$\sum_{t \in \mathbb{Z}/(NM)\mathbb{Z}} F_t(a) = (N - MD(a))F(a).$$

Cauchy–Schwarz over the NM window indices implies

$$\begin{aligned} (N - Mkh)^2 E &\leq \sum_{a \in [-h, h]^k} (N - MD(a))^2 |F(a)|^2 \\ &\leq NM \sum_t \sum_{a \in [-h, h]^k} |F_t(a)|^2 \\ &\leq NM \sum_t \sum_{a \in (\mathbb{Z}/N\mathbb{Z})^k} |F_t(a)|^2. \end{aligned} \quad (4.8)$$

The final extension is made while the summands are nonnegative squares. No restriction is discarded from a signed cube expression.

Expand each square using the difference of its two base variables. The resulting phase is $e_N(-y\tau(a))$. Apply Lemma IV.12.1 and then Lemma IV.11.1 to obtain

$$\sum_{a \in (\mathbb{Z}/N\mathbb{Z})^k} |F_t(a)|^2 \leq \prod_{\epsilon \in \{0,1\}^k} \mathcal{C}_{k+1}(\mathbf{1}_{S_t} g_\epsilon)^{1/2^k}. \quad (4.9)$$

The exponent is 2^{-k} , since the two vertices with the same first k coordinates have identical functions and phases. Hölder's inequality in t now gives

$$\frac{(N - Mkh)^2 E}{NM} \leq \prod_{\epsilon} \left(\sum_t \mathcal{C}_{k+1}(\mathbf{1}_{S_t} g_\epsilon) \right)^{1/2^k}.$$

One factor is at least the left side. Fix its ϵ . Every $t \in \mathbb{Z}/(NM)\mathbb{Z}$ occurs exactly once as $j + \ell N$, with $0 \leq j < N$ and $0 \leq \ell < M$. Averaging over j gives

$$\sum_{S \in \mathcal{P}_j} \mathcal{C}_{k+1}(\mathbf{1}_S g_\epsilon) \geq \frac{E(N - Mkh)^2}{N^2 M}$$

for some j . Dividing by MU^{k+2} and using $N - Mkh = M(w - kh)$ proves (4.6). $\square$

The same argument can be visualized with real windows $[u, u + w)$ moving around a circle of circumference N . A cube of span $D(a)$ belongs to these windows for a set of starting positions of total length $w - D(a)$. The construction in $\mathbb{Z}/(NM)\mathbb{Z}$ is an exactly finite version of that averaging argument.

IV.13.3 A uniform finite coefficient and the original input

Suppose $m = 2h + 1 > 3k$ and $m^2 \leq N$. Take

$$T = (k + 2)h, \quad M = \left\lfloor \frac{N}{T} \right\rfloor, \quad w = \frac{N}{M}. \quad (4.10)$$

Then $M \geq 4$. For $k = 1$, $m \geq 5$ gives $m^2 \geq 6(m - 1) = 4T$. For $k \geq 2$, the function $m^2 - 2(k + 2)(m - 1)$ increases in the range $m \geq 3k + 1$, and at its left endpoint equals $3k^2 - 6k + 1 \geq 1$. Hence again $N \geq 4T$. Consequently

$$T \leq w < \frac{5}{4}T, \quad U = \lceil w \rceil \leq \frac{3}{2}T, \quad w - kh \geq 2h. \quad (4.11)$$

Here $T \geq 6$ absorbs the rounding in U . Also $kh < N/6$ and $w \leq N/4$, so (4.1) holds. Every output cell has at least $T \geq h \geq m/3$ points and at most $3km$ points. By (4.7) and $m/h > 2$,

$$\mathcal{B}_k(g_\epsilon; \mathcal{P}_j, U) \geq \frac{4\alpha(m/h)^k}{(3/2)^{k+2}(k+2)^{k+2}} > \alpha \left(\frac{4}{3(k+2)} \right)^{k+2}. \quad (4.12)$$

In particular the conclusion is a *strict* failure of uniformity at the displayed coefficient.

Completion of the proof of Theorem F.21.1. Since $N \geq m^2 > 9k^2$ in this branch, $N > k + 1$ and all required factorials are invertible. To apply this to Proposition 17.7, first reduce the common difference of the input progressions to one by an invertible dilation. Write the center of the i -th progression as r_i , set $f_\epsilon(s) = f(s - \epsilon \cdot r)$, and put $\tau(a) = \sigma(a + r)$ in the normalized coordinates. These are exactly the functions in (4.5). If a dilation $x = da$ is written explicitly before this normalization, its frequency becomes $d\sigma(d(a + r))$, which is still multi-affine. After choosing ϵ , translate back using

$$g_\epsilon(x + \epsilon \cdot r) = f(x)e_N(-\phi_\epsilon(x + \epsilon \cdot r)). \quad (4.13)$$

The corresponding support partition is $\mathcal{P}_j - \epsilon \cdot r$. The phase still has degree at most $k + 1$. This identity fixes the sign of the support translation without relying on a convention in the source text.

For completeness, if $m \leq 3k$, Cauchy–Schwarz in each Fourier sum and AM–GM over its 2^k vertices give

$$E \leq Nm^k \sum_s |f(s)|^{2^{k+1}}.$$

The hypothesis therefore implies $\sum_s |f(s)|^{2^{k+1}} \geq \alpha N$. Use the singleton partition and the zero phase. With the convention $U = \ell + 1 = 2$, its quotient is at least $\alpha/2^{k+2}$, which strictly exceeds the right side of (4.12). Singletons satisfy the required lower size $m/(3k) \leq 1$. This proves the finite assertion, including the small branch, with

$$c_k = \left(\frac{4}{3(k+2)} \right)^{k+2}.$$

□

Remark IV.13.3 (Effect on subsequent uses). The new cells may be larger than those in Gowers' original construction. Their sizes remain between $m/(3k)$ and $3km$, but a later transfer to a smaller cyclic group must recheck its no-wrap condition. In the asymptotic induction, requiring $N > 2(k+1)3km$ is a sufficient safe condition and can be included in the lower threshold for N . The local phase-removal theorem itself does not require that additional inequality. No improved final Szemerédi bound is asserted here.

IV.13.4 Optimal width within this averaging argument

Proposition IV.13.4 (Width optimization). *For fixed k , as $m \rightarrow \infty$ through odd integers and $N \geq m^2$, the choice (4.10) gives*

$$\mathcal{B}_k(g; \mathcal{Q}, U) \geq (1 - o(1))\alpha \left(\frac{2}{k+2} \right)^{k+2}. \quad (4.14)$$

The real-valued model of (4.7), with $U = w$, is maximized at $w = (k+2)h$. In a subsequent limit $k \rightarrow \infty$,

$$\left(\frac{2}{k+2} \right)^{k+2} = (4e^{-2} + o(1)) \frac{2^k}{k^{k+2}}. \quad (4.15)$$

Proof. The logarithmic derivative of $(w - kh)^2/w^{k+2}$ is $2/(w - kh) - (k+2)/w$, which vanishes exactly at $(k+2)h$. It is positive before that point and negative afterwards. For fixed k , $T = (k+2)h \rightarrow \infty$ and $T/N \rightarrow 0$. Thus $w/T \rightarrow 1$, $U/w \rightarrow 1$, and $m/h \rightarrow 2$. Substitution in (4.7) gives (4.14). Finally $(1 + 2/k)^{-k-2} \rightarrow e^{-2}$, proving (4.15). The two limits have been specified in this order; no unqualified joint asymptotic is used. □

For given finite N, m, k , one can evaluate $m^k(N/M - kh)^2/(\lfloor N/M \rfloor + 1)^{k+2}$ over the admissible integer M . This gives the best bound furnished by the exact theorem, without making a continuity or rounding approximation.

IV.13.5 A geometric ceiling for every genuinely local conclusion

Let $A_d(L)$ be the number of additive d -cubes all of whose vertices belong to the integer interval $\{0, \dots, L-1\}$. For fixed increments the vertex span is the sum of their absolute values, so

$$A_d(L) = \sum_{h \in \mathbb{Z}^d} \left(L - \sum_i |h_i| \right)_+ , \quad \sum_{L \geq 0} A_d(L) z^L = \frac{z(1+z)^d}{(1-z)^{d+2}}. \quad (4.16)$$

Coefficient extraction gives

$$A_d(L) = \sum_{j=0}^d \binom{d}{j} \binom{L+d-j}{d+1} = \frac{2^d}{(d+1)!} L^{d+1} + O_d(L^d). \quad (4.17)$$

The binomial coefficient is zero when its nonnegative upper argument is less than its lower argument. The generating function follows from $\sum_{n \in \mathbb{Z}} z^{|n|} = (1+z)/(1-z)$ and $\sum_{L \geq 0} (L-r)_+ z^L = z^{r+1}/(1-z)^2$.

Proposition IV.13.5 (Local-scale upper bound). *Fix k and positive constants a_k, b_k . Suppose a phase-removal conclusion is required to use arithmetic progression cells with sizes between $a_k m$ and $b_k m$, and uses a normalization upper bound at least as large as every cell. As $m \rightarrow \infty$, $N \geq m^2$ prime, no universal coefficient in that conclusion can exceed*

$$\frac{2^{k+1}}{(k+2)!} + o(1). \quad (4.18)$$

In particular the leading logarithmic order $-k \log k$ of the coefficients obtained above is optimal within this local-scale class.

Proof. Take $f \equiv 1$ and $\sigma \equiv 0$; the input has $E = N^2 m^k$, so $\alpha = 1$. Every polynomial twist has modulus one. For $|g| \leq 1$, the triangle inequality in the expanded cube form gives

$$0 \leq \mathcal{C}_d(\mathbf{1}_Q g) \leq \mathcal{C}_d(\mathbf{1}_Q).$$

A dilation normalizes any nontrivial progression in the prime cyclic group to a consecutive interval. For $N > 2d|Q|$, its cube count agrees with the integer count: all congruences between the lifted cube vertices have difference of absolute value below N , hence are equalities. This no-wrap condition holds eventually, uniformly over the allowed cell sizes. Apply (4.17) with $d = k+1$, sum over the partition, and divide by its stated normalization. The result is at most $2^{k+1}/(k+2)! + O_k(1/m)$.

Stirling's formula gives $\log(2^{k+1}/(k+2)!) = -k \log k + O(k)$, and the coefficient in (4.14) has the same leading logarithm. Their ratio is

$$\frac{(k+2)^{k+2}}{2(k+2)!} \sim \frac{e^{k+2}}{2\sqrt{2\pi(k+2)}}.$$

Thus an exponential-in- k gap remains between the lower and upper constants; the exact optimum is not determined. $\square$




Figure IV.13.1: Negative binary logarithms of retained local factors. A smaller value means a stronger lower guarantee. The source and finite sliding curves are finite guarantees; the optimized sliding curve and the local-scale ceiling are asymptotic in the stated order of limits. The vertical axis is logarithmic.

[write] 6 October 2026; sources 15–20. The first upper bound in this report for local-scale coefficients at Proposition 17.7. It does not bind `proposition_17_7`, which has no upper cell size (the source says so after the proof).

[write] 7 October 2026; sources 60–69. *Attained asymptotically* (source 63, Section IV.67): a positive face cover retains v_k/D_k with $D_k \leq e^{2/(3(k+2))}$, so the best local coefficient is asymptotic to this ceiling (its Theorem IV.68.1 and Proposition IV.68.7). Source 63 uses this ceiling and the sliding windows of this chapter without citing this source.

This upper bound concerns conclusions that remain at a local scale. It does not restrict an unrestricted statement that can select the entire group as one cell. The interval cube formula is an elementary geometric calculation; its role here is to identify a necessary size for the localization loss.

IV.14 A refinement when the original small cells must be kept

Source 16, delivered Section 5.

The larger windows in Section IV.13 are useful when the output scale can be transported. There is also a refinement that keeps an existing equal-width cyclic partition exactly as it is.

Let its label function be $\ell(s) = \lfloor s/w \rfloor \pmod{M}$, where $w = N/M$. For a fixed increment vector a , define the normalized pattern

$$j_{a,s}(\epsilon) = \ell(s - \epsilon \cdot a) - \ell(s) \pmod{M}. \quad (5.1)$$

Lemma IV.14.1 (At most 2^k active patterns). *For fixed a , at most 2^k different patterns $j_{a,s}$ occur as s varies.*

Proof. Put $z = \{s/w\}$ and $u = a/w$. Then

$$j_{a,s}(\epsilon) = \lfloor z - \epsilon \cdot u \rfloor \pmod{M}.$$

The coordinate $\epsilon = 0$ is zero. Each of the other $2^k - 1$ coordinates changes at most once as z runs through $[0, 1)$. Hence there are at most 2^k joint values. At a breakpoint the floor function has its value immediately to the right, so no extra boundary pattern occurs. Reduction modulo M can only merge patterns. $\square$

Proposition IV.14.2 (Active count times global count). *Let L be the number of translation-normalized patterns occurring over the whole input increment box, and let R_0 be an upper bound for the number active at any one increment vector. In the localized phase-removal argument, if the phases in (3.2) are available, some candidate phase and translated copy of the original partition satisfy*

$$\sum_i \mathcal{C}_{k+1}(\mathbf{1}_{Q_i} g) \geq \frac{E}{R_0 L M}. \quad (5.2)$$

One may take $R_0 = 2^k$. The corresponding argument that uses the global count at its first Cauchy-Schwarz step gives $E/(L^2 M)$.

Proof. Group the inner Fourier sum by normalized pattern and denote the contribution of class r by $B_r(a)$. For each a , at most R_0 of these terms are nonzero, so

$$\left| \sum_r B_r(a) \right|^2 \leq R_0 \sum_r |B_r(a)|^2.$$

Keep the sum over r . Within each class, apply Cauchy-Schwarz over its M spatial translates, giving a factor M . Enlarge the increment domain while the expression is still a sum of squares. Expand it, polarize, and use (2.3). The polynomial ϕ_ϵ is independent of r .

For any fixed r, ϵ , summing over translates permutes the whole partition, so its diagonal sum is the same quantity $\sum_i \mathcal{C}_{k+1}(\mathbf{1}_{Q_i} g_\epsilon)$. Summing over r gives one factor L , and averaging over ϵ chooses the winning candidate. The resulting upper bound for E is $R_0 L M$ times that candidate's energy, proving (5.2). $\square$

[write] **6 October 2026; sources 15–20.** *Credit* (not given by the source, although source 04's text was in the repository at its pin). Lemma IV.14.1 is source 04's Lemma IV.6.5 (at most 2^k profiles at a fixed slope), and this proposition with $R_0 = 2^k$ is source 04's Lemma IV.6.6 with the first inequality of its Theorem IV.6.7, where the global count L is bounded by B_k ; the " $E/(L^2 M)$ " form is source 06's Theorem IV.8.1 and source 01's Proposition IV.3.3. The weighted variant below is new and minor.

The active count is sharp for this pointwise argument. Choose $a_i = 2^{i-1}$, a prime $N > 2^k$, and an input box containing these integers. If $M \geq 2$ and $\gcd(M, N) = 1$, all subset sums are distinct modulo N , hence their nonzero breakpoints are distinct. As s varies, $\{Ms/N\}$ visits the full N -point grid. Every one of the 2^k joint values occurs. This only proves sharpness of the active-pattern bound, not of the resulting analytic estimate.

A weighted form is sometimes useful. For positive weights λ_r , put

$$H = \max_a \sum_{\substack{r \\ \text{active at } a}} \lambda_r.$$

Weighted Cauchy-Schwarz gives the same conclusion with $R_0 L$ replaced by $H \sum_r \lambda_r^{-1}$. Thus the finite incidence pattern can be optimized without changing the partition or the phases.

IV.15 Source 36's chapter: sharp symmetry certification at Proposition 17.2 [write]

Added in the seventh write, 6 October 2026.

The next nine sections are source 36's Sections 2–10, on the global phase-extraction step around Gowers's Proposition 17.2 over vector spaces $\mathbb{F}_p^n$, where a multilinear frequency selector need not be symmetric. *Logically the chapter stands beside source 16's chapter on Proposition 17.7* (Section IV.10) and the phase-extraction material of sources 01, 04 and 06 in this Part; it is printed here because no number printed earlier may move. Source 36's introduction, with its two main theorems, is Section F.41.

What is proved. A multilinear selector $\sigma : V^k \rightarrow V^*$ whose form is not fully symmetric carries derivative-spectrum energy at most $\vartheta_{k,p} = 1 - (1 - 1/p)^k$ against every 1-bounded f , sharp for $n \geq 2$ (Theorem F.41.1); the proof goes through a rank-profile bound averaging p^{-r} over half-ranks of alternating slices (Theorem IV.18.1) and a sharp Heisenberg orthogonality inequality (Theorem IV.17.1) with its equality and stability structure (Theorem IV.20.3). For $p > k + 1$ symmetry gives a canonical polynomial and the lossless identity $\mathcal{A}_f(\sigma) = \|fe_p(-P)\|_{U^{k+1}}^{2^{k+1}}$, pointwise before averaging (Theorem F.41.2, Section IV.21); noisy, dense-domain, coset, multiaffine and prime-power variants follow (Sections IV.22–IV.23). The bilinear inequality sharpens source 31's Lemma II.106.4, and the case $k = 1$ is a vector-space analogue of source 29's symmetry step (Section II.110); source 36 cites neither. Characteristic two is excluded (Section IV.24); sources 43 and 46, placed and not yet written, treat it.

[write] **7 October 2026; sources 40–46.** Sources 43 and 46 are now printed (Section IV.25): for symmetric trilinear forms over $\mathbb{F}_2^n$ the exact energy is $2^{-r/2}$, r the rank of $T(a, a, b) + T(a, b, b)$, and source 43 extends this chapter's threshold and bilinear estimate to every finite abelian group (notes after Theorems F.41.1 and IV.17.1). Source 46 does not cite source 36; source 43 does.

Relation to the formal statements (checked at HEAD 89c1cbc38). `proposition_17_2` (Sections17_18.lean, line 48) and `proposition_17_7` (line 140) have the exact companions `proposition_17_2_holds` (Proofs17PhaseRemoval.lean, line 373) and `proposition_17_7_holds` (Proofs17LocalizedPhaseRemoval.lean, line 2941); the formal statements are on $\mathbb{Z}/N\mathbb{Z}$, and source 36 does not claim to improve their coefficient one. Nothing in this chapter is formalized.

IV.16 Normalization and the exact cube identity

Source 36, delivered Section 2.

Lemma IV.16.1 (Fourier-square expansion). *For every $g : V \rightarrow \mathbb{C}$ and $\xi \in V^*$,*

$$|\widehat{g}(\xi)|^2 = \mathbb{E}_{x,y} g(x+y) \overline{g(x)} e_p(-\xi(y)). \quad (\text{IV.16.1})$$

Consequently

$$\mathcal{A}_f(\sigma) = \mathbb{E}_{x,\mathbf{h},y} \partial_{h_1,\dots,h_k,y} f(x) e_p(-T(\mathbf{h}, y)). \quad (\text{IV.16.2})$$

Proof. Expand the square as an average over two independent points z, x , then substitute $z = x + y$. Apply the resulting identity to $g = \partial_{\mathbf{h}} f$. Commutation of multiplicative derivatives gives (IV.16.2). $\square$

For $d \geq 2$, the normalized Gowers power is

$$\|f\|_{U^d}^{2^d} = \mathbb{E}_{x, h_1, \dots, h_d} \partial_{h_1, \dots, h_d} f(x) = \mathbb{E}_{h_1, \dots, h_{d-1}} \left| \mathbb{E}_x \partial_{h_1, \dots, h_{d-1}} f(x) \right|^2. \quad (\text{IV.16.3})$$

The second identity proves the nonnegativity needed in this article. No triangle inequality for U^d is used below.

The off-by-one in the notation has mathematical content. A k -fold derivative has a linear frequency if its phase originally had degree $k + 1$. Accordingly there are $k + 1$ variables in T , the polynomial in (F.41.7) has degree at most $k + 1$, and the final norm is U^{k+1} . The following table can serve as a normalization check.

Derivative order	Selector	Total form	Extracted phase / norm
1	linear	bilinear	quadratic / U^2
2	bilinear	trilinear	cubic / U^3
3	trilinear	four-linear	quartic / U^4
k	k -linear	$(k + 1)$ -linear	degree $k + 1$ / U^{k+1}

IV.17 The sharp bilinear estimate

Source 36, delivered Section 3.

IV.17.1 Statement and a direct operator proof

Let $B : V \times V \rightarrow \mathbb{F}_p$ be bilinear. Its alternating part is

$$\Omega(h, x) = B(h, x) - B(x, h). \quad (\text{IV.17.1})$$

The rank of Ω , as a map $V \rightarrow V^*$, is even; write it as $2r$. For two functions $g, v : V \rightarrow \mathbb{C}$ define

$$\mathcal{Q}_B(g, v) = \mathbb{E}_h \left| \mathbb{E}_x g(x + h) \overline{v(x)} e_p(-B(h, x)) \right|^2, \quad \mathcal{Q}_B(g) = \mathcal{Q}_B(g, g). \quad (\text{IV.17.2})$$

The L^2 norms below use normalized counting measure.

Theorem IV.17.1 (Alternating-rank orthogonality). *For arbitrary complex functions g, v ,*

$$\mathcal{Q}_B(g, v) \leq p^{-r} \|g\|_2^2 \|v\|_2^2. \quad (\text{IV.17.3})$$

In particular $\mathcal{Q}_B(g) \leq p^{-r} \|g\|_2^4$. For each B the factor p^{-r} is attained with $g = v$ and $|g| = 1$ everywhere.

[write] **6 October 2026; sources 34–39.** This is the sharp, rank-sensitive form of source 31’s Lemma II.106.4 (note in Section F.41). Independent checks of the intake: on 21 random triples over $\mathbb{F}_3^2, \mathbb{F}_5^2, \mathbb{F}_3^3, \mathbb{F}_3^4$ the normalized ratio stayed below 0.49; the value $1/p$ is attained at $p = 3, 5, 7$; over all 6,561 trilinear forms on $\mathbb{F}_3^2$ with $f = 1$ the maximum of the tested energy over the nonsymmetric forms is exactly $5/9 = 1 - (2/3)^2$, the threshold of Theorem F.41.1 at $k = 2$.

[write] **7 October 2026; sources 40–46.** Source 43’s Theorem IV.35.2 proves the bound d_B^{-1} , $d_B = \sqrt{|G/R_B|}$, for every biadditive form on every finite abelian group, without division by two (checked at the write on small groups; note in Section F.48).

We first establish the only representation-theoretic fact required for the inequality. Inner products are linear in the first argument.

Lemma IV.17.2 (Minimum dimension of an invariant subspace). *Suppose unitary operators $(U_h)_{h \in V}$ on a finite-dimensional complex Hilbert space satisfy*

$$U_h U_l = e_p(\Omega(h, l)/2) U_{h+l}, \quad U_0 = I. \quad (\text{IV.17.4})$$

Every nonzero subspace invariant under all U_h has dimension at least $D = p^r$.

Proof. An alternating form of rank $2r$ has vectors $a_1, \dots, a_r, b_1, \dots, b_r$ with $\Omega(a_i, b_j) = \delta_{ij}$ and all pairings within the a -family or within the b -family zero. To obtain them, select a, b with $\Omega(a, b) \neq 0$, scale b to make the value one, split off their nondegenerate two-dimensional span, and repeat on its orthogonal complement. This also proves evenness of the rank.

Put $X_i = U_{a_i}$ and $Z_i = U_{b_i}$. The Z_i commute, $Z_i^p = I$, and

$$X_i Z_j = e_p(\delta_{ij}) Z_j X_i.$$

Let E be a nonzero invariant subspace. It is invariant under the inverses as well, because those are among the U_h up to scalars. The commuting unitary restrictions $Z_i|_E$ have a common nonzero eigenvector z . For each $t = (t_1, \dots, t_r) \in \mathbb{F}_p^r$, the vector $X_1^{t_1} \cdots X_r^{t_r} z$ belongs to E . Its simultaneous Z_i eigenvalues are those of z multiplied by $e_p(-t_i)$. Different t give different joint eigenspaces, which are orthogonal. There are p^r such nonzero vectors. The case $r = 0$ is immediate. $\square$

Proof of the inequality in Theorem IV.17.1. Work on $\mathcal{H} = L^2(V)$ and put

$$W_h g(x) = g(x + h) e_p(-B(h, x)), \quad (\text{IV.17.5})$$

$$U_h = e_p(-B(h, h)/2) W_h. \quad (\text{IV.17.6})$$

A direct multiplication gives $W_h W_l = e_p(-B(l, h)) W_{h+l}$; substituting (IV.17.6) yields (IV.17.4). For $u \in \mathcal{H}$ let P_u be the positive rank-one operator $P_u(z) = \langle z, u \rangle u$. Its trace is $\|u\|_2^2$. Define the positive operator

$$M_g = \mathbb{E}_h P_{U_h g}. \quad (\text{IV.17.7})$$

Translation of the index h , together with the cocycle law, gives $U_l M_g U_l^{-1} = M_g$ for all l . Each eigenspace of M_g is therefore invariant. If $\lambda > 0$ is its largest eigenvalue, the corresponding eigenspace has dimension at least D by Lemma IV.17.2. Thus

$$D\lambda \leq \text{tr}(M_g) = \|g\|_2^2.$$

Finally, multiplication by the scalar in (IV.17.6) does not change the absolute value of a matrix coefficient, and

$$\mathcal{Q}_B(g, v) = \langle M_g v, v \rangle \leq \lambda \|v\|_2^2.$$

If $M_g = 0$ the conclusion is automatic. This proves (IV.17.3) without an irreducible-representation classification. $\square$

IV.17.2 Sharp unit-modulus examples

The symplectic splitting in the preceding proof yields independent functionals $\alpha_1, \dots, \alpha_r, \beta_1, \dots, \beta_r \in V^*$ such that

$$\Omega(h, x) = \sum_{j=1}^r (\alpha_j(h) \beta_j(x) - \beta_j(h) \alpha_j(x)). \quad (\text{IV.17.8})$$

Set $B_0(h, x) = \sum_j \alpha_j(h) \beta_j(x)$. Then $S = B - B_0$ is symmetric. With $P_B(x) = S(x, x)/2$ and $g(x) = e_p(P_B(x))$, direct expansion gives

$$\mathbb{E}_x \partial_h g(x) e_p(-B(h, x)) = e_p(S(h, h)/2) \mathbb{E}_x e_p(-B_0(h, x)). \quad (\text{IV.17.9})$$

The last average is one when all $\alpha_j(h)$ vanish and zero otherwise, by character orthogonality. That event has probability p^{-r} . This proves the sharpness assertion of Theorem IV.17.1, including the 1-bounded setting.

For instance, on $\mathbb{F}_p^2$ take $B(h, x) = h_1 x_2$ and $g = 1$. The alternating rank is two, and $\mathcal{Q}_B(g) = 1/p$. On $\mathbb{F}_p^{2r}$, taking $B(h, x) = \sum_{j=1}^r h_{2j-1} x_{2j}$ and $g = 1$ gives p^{-r} . A bound based only on whether the alternating part vanishes would lose all information distinguishing these examples.

IV.18 The derivative rank profile

Source 36, delivered Section 4.

Fix $i \in \{1, \dots, k\}$. Let $\mathbf{u}$ denote the other $k-1$ increments, inserted in their original slots, and set

$$g_{i,\mathbf{u}}(x) = \partial_{h_1, \dots, \widehat{h_i}, \dots, h_k} f(x), \quad m_i(\mathbf{u}) = \|g_{i,\mathbf{u}}\|_2^2, \quad (\text{IV.18.1})$$

$$B_{i,\mathbf{u}}(a, b) = T(h_1, \dots, h_{i-1}, a, h_{i+1}, \dots, h_k, b), \quad (\text{IV.18.2})$$

$$\Omega_{i,\mathbf{u}}(a, b) = B_{i,\mathbf{u}}(a, b) - B_{i,\mathbf{u}}(b, a), \quad r_i(\mathbf{u}) = \frac{1}{2} \text{rank } \Omega_{i,\mathbf{u}}. \quad (\text{IV.18.3})$$

For $k = 1$ the tuple $\mathbf{u}$ is empty, has probability one, and $g_{i,\mathbf{u}} = f$. These conventions avoid a separate base case later.

Theorem IV.18.1 (Amplitude-sensitive rank-profile bound). *For every multilinear selector and every complex f ,*

$$\mathcal{A}_f(\sigma) \leq \min_{1 \leq i \leq k} \mathbb{E}_{\mathbf{u}} p^{-r_i(\mathbf{u})} m_i(\mathbf{u})^2. \quad (\text{IV.18.4})$$

For 1-bounded f this implies

$$\mathcal{A}_f(\sigma) \leq \min_i \mathbb{E}_{\mathbf{u}} p^{-r_i(\mathbf{u})}. \quad (\text{IV.18.5})$$

Proof. Fix i and average last over $h_i = a$. Commutation of derivatives gives $\partial_h f = \partial_a g_{i,\mathbf{u}}$. Therefore

$$\mathcal{A}_f(\sigma) = \mathbb{E}_{\mathbf{u}} \mathcal{Q}_{B_{i,\mathbf{u}}}(g_{i,\mathbf{u}}).$$

Apply Theorem IV.17.1 to each slice. If $|f| \leq 1$, then $|g_{i,\mathbf{u}}| \leq 1$ and $m_i(\mathbf{u}) \leq 1$. Since the calculation applies to every i , both minima follow. $\square$

Corollary IV.18.2 (Quantitative low-rank consequences). *Assume $|f| \leq 1$ and $\mathcal{A}_f(\sigma) \geq a$. For every i and integer $t \geq 1$,*

$$\mathbb{P}_{\mathbf{u}}(r_i(\mathbf{u}) \geq t) \leq \frac{1-a}{1-p^{-t}}, \quad (\text{IV.18.6})$$

$$\mathbb{P}_{\mathbf{u}}(r_i(\mathbf{u}) < t) \geq \frac{a-p^{-t}}{1-p^{-t}}. \quad (\text{IV.18.7})$$

The second bound is informative when $a > p^{-t}$. If $r_i \geq t$ on a set of tuples of density at least ρ , then

$$\mathcal{A}_f(\sigma) \leq 1 - \rho(1 - p^{-t}). \quad (\text{IV.18.8})$$

Proof. The profile average is at most $\mathbb{P}(r_i < t) + p^{-t}\mathbb{P}(r_i \geq t)$. Rearrange and use Theorem IV.18.1. $\square$

The amplitude information should not be discarded unnecessarily. If a frozen derivative has small L^2 mass, it is expensive to attribute energy to that slice even when its alternating rank is zero. Conversely, a large unweighted set of rank-zero slices need not carry appreciable energy. Formula (IV.18.4) keeps these distinctions.

As a simple scale check, for $p = 101$, $k = 2$, an energy of 0.05 is above $\vartheta_{2,101} = 201/10201$. It therefore certifies symmetry. For $p = 5$ the same energy is below $\vartheta_{2,5} = 9/25$ and does not do so. The rank profile may still give useful information in the latter case.

IV.19 Proof and sharpness of the symmetry threshold

Source 36, delivered Section 5.

IV.19.1 A support lemma

Lemma IV.19.1 (Nonzero multilinear maps have positive support). *Let $F : V_1 \times \cdots \times V_m \rightarrow W$ be a nonzero multilinear map of finite-dimensional $\mathbb{F}_p$ -vector spaces. Then*

$$\mathbb{P}(F \neq 0) \geq (1 - 1/p)^m. \quad (\text{IV.19.1})$$

The convention for $m = 0$ is a nonzero constant, with support probability one. The same bound holds when F is separately affine.

Proof. For a nonzero linear map in one variable, its kernel has codimension at least one, so its zero probability is at most $1/p$. A nonzero affine map either has no zeros or has a zero set which is a translate of a proper linear kernel, giving the same bound.

For the induction, regard F as a map in its first $m - 1$ variables with values in the vector space of linear maps $V_m \rightarrow W$ (or the vector space of affine maps in the separately affine case). This new map is nonzero and of the same separate type. By induction, it takes a nonzero value on at least a fraction $(1 - 1/p)^{m-1}$ of tuples. Conditional on such a tuple, the remaining nonzero linear or affine map is nonzero at a fraction at least $1 - 1/p$ of its inputs. Multiplying the two fractions proves the assertion. $\square$

The bound is attained by $F(v_1, \dots, v_m) = \ell_1(v_1) \cdots \ell_m(v_m)w$, where the ℓ_j are nonzero functionals and $w \neq 0$. Vector-valued output is important in the application: the output will be a whole alternating bilinear form, not its value at a preselected pair.

IV.19.2 Certification

Proof of the inequality in Theorem F.41.1. The transpositions $(i, k + 1)$, $1 \leq i \leq k$, generate the symmetric group on the $k + 1$ slots. Therefore, if T is not fully symmetric, there is some i for which

$$T(h_1, \dots, h_i, \dots, h_k, y) - T(h_1, \dots, y, \dots, h_k, h_i)$$

is nonzero as a multilinear form. In the notation of (IV.18.3), this says that $\mathbf{u} \mapsto \Omega_{i,\mathbf{u}}$ is a nonzero $(k-1)$ -linear map into the vector space of alternating bilinear forms on V . Lemma IV.19.1 gives

$$\rho := \mathbb{P}_{\mathbf{u}}(\Omega_{i,\mathbf{u}} \neq 0) \geq (1 - 1/p)^{k-1}.$$

Every nonzero alternating form has rank at least two, so $r_i(\mathbf{u}) \geq 1$ on this event. Hence

$$\begin{aligned} \mathcal{A}_f(\sigma) &\leq \mathbb{E}_{\mathbf{u}} p^{-r_i(\mathbf{u})} \\ &\leq (1 - \rho) + \rho/p \\ &\leq 1 - (1 - 1/p)^k. \end{aligned}$$

This includes $k = 1$ because the nonzero defect is then a fixed alternating bilinear form. $\square$

The proof combines two independent exact obstructions: a nonzero multilinear defect cannot hide on too few slices, and on every defective slice noncommutation forces a factor at most $1/p$. Treating the alternating defect as an arbitrary scalar polynomial would obscure the second obstruction and lose the sharp constant.

IV.19.3 Endpoint construction

Proof of sharpness in Theorem F.41.1. Assume $\dim V \geq 2$ and choose linearly independent $a, c \in V^*$. Set $f = 1$ and

$$\sigma(h_1, \dots, h_k) = \left(\prod_{j=1}^k a(h_j) \right) c. \quad (\text{IV.19.2})$$

This selector is symmetric in its k inputs. Choose $u, v \in V$ with $a(u) = 1$, $c(u) = 0$, $a(v) = 0$, $c(v) = 1$. Then $T(u, \dots, u, v) = 1$, whereas swapping the first and last entries makes the value zero. Thus T is not fully symmetric.

Every multiplicative derivative of f is one. Its Fourier coefficient at a frequency is one exactly when that frequency is zero. Therefore

$$\mathcal{A}_1(\sigma) = \mathbb{P} \left(\prod_{j=1}^k a(h_j) = 0 \right) = 1 - (1 - 1/p)^k. \quad (\text{IV.19.3})$$

For a fixed i , the alternating slice is a scalar multiple of $a \otimes c - c \otimes a$. Its rank is zero when $\prod_{j \neq i} a(h_j) = 0$ and two otherwise. Thus the rank-profile bound and the support lemma are simultaneously equalities. $\square$

Remark IV.19.2 (The strict inequality cannot be relaxed). The implication in Theorem F.41.1 uses $>$, not $\geq$. It remains sharp after imposing symmetry in the derivative slots or restricting to unit-modulus functions. In dimension one every scalar multilinear form is already fully symmetric, so the nonsymmetric extremal class is empty there.

For reference, the first three thresholds are as follows. These are symmetry thresholds; classical integration separately requires $p > k + 1$.

p	$k = 1$	$k = 2$	$k = 3$
3	1/3	5/9	19/27
5	1/5	9/25	61/125
17	1/17	33/289	817/4913
101	1/101	201/10201	30301/1030301

IV.20 Equality structure and stability of the bilinear step

Source 36, delivered Section 6.

The short operator proof suffices for the main threshold. A fuller orthogonality identity explains how close-to-extremal slices look. This section is a refinement of the bilinear step, not a classification of all extremizers of the higher-order threshold.

IV.20.1 Block decomposition and an exact identity

Retain B , Ω , r , $D = p^r$, and the operators U_h from Section IV.17. Put $R = \text{rad } \Omega$ and choose a nondegenerate complement W , so $V = R \oplus W$ and $|W| = D^2$. The U_z , $z \in R$, form an ordinary commuting representation of the additive group R . Consequently

$$\mathcal{H} = \bigoplus_{\chi \in R_{\text{char}}^*} \mathcal{H}_\chi, \quad U_z|_{\mathcal{H}_\chi} = \chi(z)I. \quad (\text{IV.20.1})$$

Here R_{char}^* denotes its complex character group. Zero blocks may be omitted for a general representation; for the concrete representation on $L^2(V)$ every block below is nonzero.

Lemma IV.20.1 (Explicit Weyl decomposition). *There are unitary identifications*

$$\mathcal{H}_\chi \simeq \mathbb{C}^D \otimes \mathbb{C}^D \quad (\text{IV.20.2})$$

under which U_w , $w \in W$, acts on the first factor and trivially on the second. Writing its first-factor action as $A_{\chi,w}$, the D^2 matrices $D^{-1/2}A_{\chi,w}$ form an orthonormal basis of $\text{Mat}_D(\mathbb{C})$ for the Hilbert–Schmidt inner product.

Proof. Choose symplectic generators X_i, Z_i as in Lemma IV.17.2. Within a fixed block, the common Z_i eigenspaces are permuted transitively by the X_i , so they have a common dimension, say m_χ . Choose an orthonormal basis of the joint eigenspace with all eigenvalues one, and transport it by the X_i . That eigenspace exists whenever the block is nonzero, because the X_i shift any joint eigenvalue tuple to $(1, \dots, 1)$. In the resulting basis the X_i shift the D eigenvalue labels and the Z_i multiply them by characters, without affecting the basis index within the original eigenspace. This constructs $\mathcal{H}_\chi \simeq \mathbb{C}^D \otimes \mathbb{C}^{m_\chi}$.

For the concrete U_h on $L^2(V)$, the trace of U_h is zero for $h \neq 0$: it is a scalar times a modulated nonzero translation and has zero diagonal in the point-mass basis. The projector onto $\mathcal{H}_\chi$ is $|R|^{-1} \sum_{z \in R} \overline{\chi(z)} U_z$. Its trace is $|V|/|R| = D^2$. Thus $m_\chi = D$ for every χ .

If $w \neq 0$ lies in W , nondegeneracy provides $u \in W$ with $\Omega(u, w) \neq 0$. Conjugation by $A_{\chi,u}$ multiplies $A_{\chi,w}$ by $e_p(\Omega(u, w)) \neq 1$, so its trace is zero. The cocycle law then gives $\text{tr}(A_{\chi,w} A_{\chi,w'}^*) = 0$ for $w \neq w'$, and equals D for $w = w'$. There are D^2 orthogonal matrices in the D^2 -dimensional matrix space, proving the last assertion. $\square$

Identify a vector in each tensor block with its coefficient matrix. Write G_χ and V_χ for the matrices of the projections of g and v . Their Hilbert–Schmidt norms equal the Hilbert norms of those projections.

Proposition IV.20.2 (Exact matrix-coefficient formula). *With the preceding identifications,*

$$\mathcal{Q}_B(g, v) = \frac{1}{D} \sum_\chi \|G_\chi V_\chi^*\|_{\text{HS}}^2. \quad (\text{IV.20.3})$$

In particular, with $\rho_\chi = G_\chi G_\chi^*$,

$$D\mathcal{Q}_B(g) = \sum_\chi \text{tr}(\rho_\chi^2), \quad \sum_\chi \text{tr} \rho_\chi = \|g\|_2^2. \quad (\text{IV.20.4})$$

Proof. Write $h = z + w$ with $z \in R$, $w \in W$. Since $\Omega(z, w) = 0$, $U_h = U_z U_w$. The matrix coefficient in question is

$$\sum_\chi \chi(z) \text{tr}(G_\chi V_\chi^* A_{\chi, w}).$$

Averaging its squared modulus over z eliminates all cross terms between different χ . By the orthonormal-basis assertion in Lemma IV.20.1, for every $D \times D$ matrix K ,

$$\mathbb{E}_{w \in W} |\text{tr}(K A_{\chi, w})|^2 = D^{-1} \|K\|_{\text{HS}}^2.$$

Averaging over w proves (IV.20.3). Set $v = g$ to obtain (IV.20.4). $\square$

The mixed inequality also follows from $\|GV^*\|_{\text{HS}} \leq \|G\|_{\text{HS}} \|V\|_{\text{HS}}$ and $\sum_\chi a_\chi b_\chi \leq (\sum_\chi a_\chi)(\sum_\chi b_\chi)$ for nonnegative sequences. This gives a second proof of its exact normalization.

IV.20.2 Rigidity and a dimension-free stability estimate

Theorem IV.20.3 (Bilinear equality and stability). *Let $\|g\|_2 = 1$. Equality $\mathcal{Q}_B(g) = D^{-1}$ holds if and only if g is supported in one block $\mathcal{H}_\chi$ and is a simple tensor there. Equivalently, exactly one G_χ is nonzero and it has rank one.*

If $0 \leq \varepsilon \leq 1$ and $\mathcal{Q}_B(g) \geq (1 - \varepsilon)/D$, there is a unit vector g_0 with $\mathcal{Q}_B(g_0) = 1/D$ such that

$$\|g - g_0\|_2^2 \leq 2(1 - \sqrt{1 - \varepsilon}). \quad (\text{IV.20.5})$$

When $0 \leq \varepsilon < 1/2$, the right-hand side can be improved to

$$2 \left(1 - \sqrt{\frac{1 + \sqrt{1 - 2\varepsilon}}{2}} \right). \quad (\text{IV.20.6})$$

These are assertions in $L^2(V)$; g_0 is not asserted to be 1-bounded.

Proof. Collect all eigenvalues of all the positive matrices ρ_χ into one finite list (λ_j) . Their sum is one, and $D\mathcal{Q}_B(g) = \sum_j \lambda_j^2$. This sum is one precisely when one λ_j is one and the others are zero. This proves the equality classification.

Let λ be the largest eigenvalue in the list. Since $\sum_j \lambda_j^2 \leq \lambda \sum_j \lambda_j = \lambda$, we have $\lambda \geq 1 - \varepsilon$. A singular value decomposition of the corresponding G_χ supplies a unit rank-one tensor g_0 whose inner product with g has modulus $\sqrt{\lambda}$. Choose its scalar phase to make that inner product nonnegative. Then $\|g - g_0\|_2^2 = 2(1 - \sqrt{\lambda})$, proving (IV.20.5).

If $\varepsilon < 1/2$, the same argument shows $\lambda > 1/2$. All other eigenvalues sum to $1 - \lambda$, whence

$$1 - \varepsilon \leq \sum_j \lambda_j^2 \leq \lambda^2 + (1 - \lambda)^2.$$

On $[1/2, 1]$ this implies $\lambda \geq (1 + \sqrt{1 - 2\varepsilon})/2$ and yields (IV.20.6). $\square$

The last scalar bound is sharp when two orthogonal Schmidt components are available: take their squared coefficients to be λ and $1 - \lambda$. That is a statement about the purity argument, not a claim that every such example obeys a pointwise bound. For higher-order derivative spectra, the simple tensors arising in different frozen slices must also be compatible with one common f . That compatibility is not supplied by this stability theorem.

IV.21 Lossless polarization and the Section 17 interface

Source 36, delivered Section 7.

IV.21.1 A self-contained polarization calculation

For an additive-valued function P , write $\Delta_h P(x) = P(x + h) - P(x)$, using a capital delta to distinguish it from the multiplicative derivative.

Lemma IV.21.1 (Diagonal integration). *Let $d \geq 2$, $p > d$, and let $T : V^d \rightarrow \mathbb{F}_p$ be symmetric and d -linear. If $P(x) = T(x, \dots, x)/d!$, then*

$$\Delta_{h_1} \cdots \Delta_{h_d} P(x) = T(h_1, \dots, h_d). \quad (\text{IV.21.1})$$

Moreover, with $d = k + 1$,

$$\Delta_{h_1} \cdots \Delta_{h_k} P(x) = T(h_1, \dots, h_k, x) + C(\mathbf{h}), \quad C(\mathbf{h}) = \Delta_{\mathbf{h}} P(0). \quad (\text{IV.21.2})$$

Proof. Expand the d -fold difference as an alternating sum over subsets $S \subseteq \{1, \dots, d\}$ of $P(x + \sum_{j \in S} h_j)$. Multilinearity expands each term into ordered choices of d inputs from $x, h_1, \dots, h_d$. A term missing some h_j cancels upon summing over whether j is in S . Thus only terms containing every h_j survive. There are $d!$ of them, they contain no x , and symmetry makes all of them equal to $T(h_1, \dots, h_d)$. Division by $d!$ gives (IV.21.1).

For the second assertion, set $F(x) = \Delta_{\mathbf{h}} P(x)$. The first assertion says $F(x + y) - F(x) = T(\mathbf{h}, y)$. Taking $x = 0$ gives $F(y) = T(\mathbf{h}, y) + F(0)$, which is (IV.21.2). $\square$

Proof of Theorem F.41.2. Multiplicative derivatives distribute over products. Therefore, by (IV.21.2),

$$\begin{aligned} \partial_{\mathbf{h}}(f e_p(-P))(x) &= \partial_{\mathbf{h}} f(x) e_p(-\Delta_{\mathbf{h}} P(x)) \\ &= e_p(-C(\mathbf{h})) \partial_{\mathbf{h}} f(x) e_p(-\sigma(\mathbf{h})(x)). \end{aligned} \quad (\text{IV.21.3})$$

Average over x and take absolute values. This is (F.41.8). Squaring and averaging over $\mathbf{h}$, then using (IV.16.3), gives (F.41.9). $\square$

The proof does not use a choice among 2^{k+1} candidate phases or introduce an additional Cauchy–Schwarz power. In a homogeneous global model the phase is read directly from the diagonal of T . One may add any polynomial of degree at most k to P without changing the pointwise magnitudes: its k -fold derivative is constant in x . Thus “canonical” denotes the specified homogeneous representative, not uniqueness among all possible phases.

IV.21.2 Exact scalar specialization

Corollary IV.21.2 (Equality-form replacement for the global scalar step). *Let $V = \mathbb{F}_p$, $p > k + 1$, and $\sigma(h_1, \dots, h_k) = ch_1 \cdots h_k$. For every $f : V \rightarrow \mathbb{C}$, with $P(x) = cx^{k+1}/(k+1)!$ and $g = fe_p(-P)$,*

$$\sum_{\mathbf{h} \in V^k} \left| \sum_{x \in V} \partial_{\mathbf{h}} f(x) e_p(-ch_1 \cdots h_k x) \right|^2 = \sum_{\mathbf{h} \in V^k} \left| \sum_{x \in V} \partial_{\mathbf{h}} g(x) \right|^2. \quad (\text{IV.21.4})$$

Proof. The total form is $ch_1 \cdots h_k y$, hence is fully symmetric. Apply (F.41.8) and multiply both sides after averaging by p^{k+2} . $\square$

This is the scale N^{k+2} used in Proposition 17.2. The original paper uses backward cubes

$$\prod_{\epsilon \in \{0,1\}^k} \mathcal{C}^{|\epsilon|} f(s - \epsilon \cdot \mathbf{h}),$$

where $\mathcal{C}$ denotes conjugation. Our forward cube is its translate by $h_1 + \cdots + h_k$; its Fourier coefficient differs only by a unit scalar. Thus the squared quantities in (IV.21.4) match the source convention. The corrected repository already raises the phase degree to $k + 1$ and includes the characteristic restriction; these are not newly reported defects here [166].

Example IV.21.3 (Why the degree must be $k + 1$). Take $k = 1$, $V = \mathbb{F}_p$, $a \neq 0$, $f(x) = e_p(ax^2)$, and $\sigma(h)(x) = 2ahx$. Every tested Fourier coefficient has modulus one, so $\mathcal{A}_f(\sigma) = 1$. Twisting by any linear polynomial preserves the quadratic coefficient. For such a twist g , $\mathbb{E}_x \partial_h g(x)$ vanishes for $h \neq 0$ and has modulus one for $h = 0$. Hence $\|g\|_{U^2}^4 = 1/p$, not one. The quadratic phase in Theorem F.41.2 is necessary.

IV.21.3 An elementary complete quadratic inverse consequence

Corollary IV.21.4. *Suppose $p > 2$, $k = 1$, $|f| \leq 1$, and $\mathcal{A}_f(\sigma) = a > 1/p$. There is a polynomial Q of degree at most two such that*

$$|\mathbb{E}_x f(x) e_p(-Q(x))| \geq \sqrt{a / \|f\|_2^2} \geq \sqrt{a}. \quad (\text{IV.21.5})$$

Proof. Theorems F.41.1 and F.41.2 give a quadratic P with $g = fe_p(-P)$ and $\|g\|_{U^2}^4 = a$. Finite character orthogonality, applied to the Fourier expansion of g , gives Parseval and the fourth-moment identity

$$\sum_{\xi \in V^*} |\widehat{g}(\xi)|^2 = \|g\|_2^2, \quad \sum_{\xi \in V^*} |\widehat{g}(\xi)|^4 = \|g\|_{U^2}^4.$$

For completeness, in the latter identity substitute $g(x) = \sum_{\xi} \widehat{g}(\xi) e_p(\xi(x))$ into the parallelogram average; orthogonality in its three variables forces the four frequencies to agree. Hence $a \leq \max_{\xi} |\widehat{g}(\xi)|^2 \|g\|_2^2$. Choose a maximizing ξ and put $Q = P + \xi$. $\square$

For larger k , the lossless identity supplies an input to a U^{k+1} -inverse theorem. It does not itself prove the required lower-order inverse theorem or its quantitative bounds.

IV.22 Robust and local variants

Source 36, delivered Section 8.

IV.22.1 Errors in a frequency selector

The expression $\mathcal{A}_f(\phi)$ in (F.41.3) makes sense for an arbitrary map $\phi : V^k \rightarrow V^*$, whether or not it is multilinear.

Proposition IV.22.1 (Selector disagreement costs at most its density). *If $|f| \leq 1$ and $\phi, \sigma : V^k \rightarrow V^*$ disagree on a set of density at most η , then*

$$|\mathcal{A}_f(\phi) - \mathcal{A}_f(\sigma)| \leq \eta. \quad (\text{IV.22.1})$$

If σ is multilinear, $\mathcal{A}_f(\phi) \geq a$, and $a - \eta > \vartheta_{k,p}$, its total form is symmetric. If also $p > k + 1$, the canonical phase satisfies

$$\|fe_p(-P)\|_{U^{k+1}}^{2^{k+1}} \geq a - \eta. \quad (\text{IV.22.2})$$

Proof. On agreement tuples the summands coincide. On each remaining tuple, both squared magnitudes lie in $[0, 1]$, so their difference in absolute value is at most one. Average, then apply the main theorems. $\square$

This is an agreement assumption on the selectors, not a claim that a multilinear fit can be found efficiently from arbitrary noisy data. Obtaining the fit remains a separate structural or algorithmic problem. The inequality also allows deterministic rather than random corruption.

IV.22.2 A dense set of large coefficients

Corollary IV.22.2 (Dense-domain extraction). *Let $H \subseteq V^k$ have density β , let σ be a global multilinear selector, and suppose $|\widehat{\partial_{\mathbf{h}} f}(\sigma(\mathbf{h}))| \geq \gamma$ for every $\mathbf{h} \in H$. If $|f| \leq 1$, then*

$$\mathcal{A}_f(\sigma) \geq \beta\gamma^2. \quad (\text{IV.22.3})$$

If $\beta\gamma^2 > \vartheta_{k,p}$, the form T is symmetric. If additionally $p > k + 1$, its canonical phase obeys

$$\|fe_p(-P)\|_{U^{k+1}} \geq (\beta\gamma^2)^{1/2^{k+1}}. \quad (\text{IV.22.4})$$

A sufficient characteristic condition for both conclusions is $p > \max\{k + 1, k/(\beta\gamma^2)\}$.

Proof. Retain only the nonnegative summands indexed by H . The characteristic condition follows from (F.41.6). $\square$

It is essential that σ is defined multilinearly on all of V^k . Values on H alone do not supply such an extension. In particular, a local multilinear pattern on an arithmetic box in the cyclic-group argument cannot be inserted into this corollary without an additional extension theorem.

IV.22.3 Common selectors on affine cosets

Let $W \leq V$, choose a set $\mathcal{C}$ of representatives of affine cosets of W , and assign nonnegative weights $\mu(c)$ of total one. For $c \in \mathcal{C}$, put $f_c(w) = f(c + w)$. Let a single multilinear $\sigma : W^k \rightarrow W^*$ be used in every coset and define

$$\mathcal{A}_{\mu, f}(\sigma) = \sum_{c \in \mathcal{C}} \mu(c) \mathcal{A}_{f_c}(\sigma). \quad (\text{IV.22.5})$$

Proposition IV.22.3 (Weighted coset version). *If $|f| \leq 1$ and $\mathcal{A}_{\mu,f}(\sigma) > \vartheta_{k,p}$, the total form of σ on W is symmetric. If $p > k + 1$ and P is the canonical polynomial on W , then*

$$\mathcal{A}_{\mu,f}(\sigma) = \sum_c \mu(c) \|f c e_p(-P)\|_{U^{k+1}(W)}^{2^{k+1}}. \quad (\text{IV.22.6})$$

The weighted rank-profile estimate holds with $m_i(\mathbf{u})^2$ replaced by $\sum_c \mu(c) \|\partial_{\mathbf{u}} f c\|_2^4$.

Proof. If the common total form were nonsymmetric, each coset's energy would be at most $\vartheta_{k,p}$ by Theorem F.41.1; average the inequalities. The identity and the profile assertion follow by applying their corresponding theorems separately on W and summing. $\square$

These are local norms on a subspace with its uniform counting measure. The proposition does not assert that the resulting piecewise function is a low-degree polynomial on the full ambient space. Such a gluing statement requires additional compatibility between cosets.

IV.23 Multiaffine selectors and extension fields

Source 36, delivered Section 9.

IV.23.1 Multiaffine frequency patterns

A separately affine selector has a unique decomposition

$$\sigma(\mathbf{h}) = \sum_{S \subseteq [k]} \sigma_S((h_j)_{j \in S}), \quad (\text{IV.23.1})$$

where σ_S is multilinear in the indicated inputs and $\sigma_\emptyset$ is constant. One can obtain this decomposition by successive subtraction of the values at zero in each slot; the same procedure proves uniqueness.

For fixed i and $\mathbf{u}$, write

$$\sigma(h_1, \dots, a, \dots, h_k) = L_{i,\mathbf{u}}(a) + \xi_{i,\mathbf{u}}. \quad (\text{IV.23.2})$$

Define $B_{i,\mathbf{u}}(a, b) = L_{i,\mathbf{u}}(a)(b)$ and its alternating rank as before. These ranks refer to the linear part in the unfrozen input, not to the affine offset.

Theorem IV.23.1 (Multiaffine coefficient certification). *For a separately affine selector and arbitrary complex f , the amplitude-sensitive estimate (IV.18.4) still holds with the ranks just defined. If $|f| \leq 1$ and $\mathcal{A}_f(\sigma) > \vartheta_{k,p}$, then for every nonempty $S \subseteq [k]$ the form*

$$((h_j)_{j \in S}, y) \mapsto \sigma_S((h_j)_{j \in S})(y) \quad (\text{IV.23.3})$$

is fully symmetric in its $|S| + 1$ variables. No restriction on the constant component is asserted.

Proof. For a fixed slice set $g = \partial_{\mathbf{u}} f$ and $v(x) = e_p(\xi_{i,\mathbf{u}}(x))g(x)$. The slice energy is exactly $\mathcal{Q}_{B_{i,\mathbf{u}}}(g, v)$, and $\|v\|_2 = \|g\|_2$. The mixed form of Theorem IV.17.1 proves the profile.

The map $\mathbf{u} \mapsto \Omega_{i,\mathbf{u}}$ is separately affine into the space of alternating bilinear forms. If it is nonzero, the affine version of Lemma IV.19.1 and the rank profile imply $\mathcal{A}_f(\sigma) \leq \vartheta_{k,p}$ exactly as before. Thus under the strict energy hypothesis every such map is zero.

Expand it using (IV.23.1). For fixed i its component indexed by $S \setminus \{i\}$, with $i \in S$, is precisely the alternating defect obtained by swapping h_i and y in (IV.23.3). Uniqueness of the decomposition forces each component defect to vanish. For any fixed nonempty S , the swaps between its inputs and y generate all permutations of these $|S| + 1$ slots. The component form is therefore fully symmetric. $\square$

The homogeneous top part can consequently be integrated when $p > k + 1$. One must not conclude that its diagonal polynomial removes the lower-order components of the frequency selector. The following example shows a failure of the naive equality-form extension.

Example IV.23.2 (An affine offset survives). Let $V = \mathbb{F}_{17}$, $k = 1$, $\xi(x) = x$, and $f(x) = (1 + e_{17}(x))/2$. Take the constant selector $\sigma(h) = \xi$. It is separately affine, its linear part is zero, and

$$\partial_h f(x) = \frac{1}{4}(1 + e_{17}(h) + e_{17}(x + h) + e_{17}(-x)).$$

Thus $\mathcal{A}_f(\sigma) = 1/16 > 1/17$, although $\|f\|_{U^2}^4 = 1/8$. Integrating the zero top part gives $P = 0$, so $\mathcal{A}_f(\sigma) = \|f e_{17}(-P)\|_{U^2}^4$ would be false. There is no contradiction with Theorem IV.23.1, which certifies the symmetry of nonconstant components, not this extra identity.

IV.23.2 Finite extensions of the prime field

Corollary IV.23.3 (Odd prime-power fields). *Let $q = p^e$ with p odd, let $V = \mathbb{F}_q^n$, and let $\psi : \mathbb{F}_q \rightarrow \mathbb{C}^\times$ be a nontrivial additive character. Use $\mathbb{F}_q$ -linear duals and replace e_p by ψ in the definitions. The rank-profile bound becomes*

$$\mathcal{A}_f(\sigma) \leq \min_i \mathbb{E}_{\mathbf{u}} q^{-r_i(\mathbf{u})} m_i(\mathbf{u})^2, \quad (\text{IV.23.4})$$

where ranks are over $\mathbb{F}_q$. The sharp symmetry threshold is $1 - (1 - 1/q)^k$. If the total form is symmetric and the characteristic satisfies $p > k + 1$, the canonical polarization and lossless identities hold over $\mathbb{F}_q$ as well.

Proof. Write $\psi(t) = e_p(\text{Tr}_{\mathbb{F}_q/\mathbb{F}_p}(ct))$ for some $c \neq 0$. A bilinear alternating form Ω of $\mathbb{F}_q$ -rank $2r$ induces the $\mathbb{F}_p$ -bilinear alternating form $\text{Tr}(c\Omega)$ of rank $2er$. Indeed the radicals are the same sets: if h is outside the $\mathbb{F}_q$ -radical, the $\mathbb{F}_q$ -linear map $x \mapsto \Omega(h, x)$ is onto $\mathbb{F}_q$, and the nonzero trace functional does not vanish on its entire image. The dimension of the common radical is multiplied by e when scalars are restricted. Theorem IV.17.1, on the underlying $\mathbb{F}_p$ -space, therefore gives the factor $p^{-er} = q^{-r}$.

A nonzero $\mathbb{F}_q$ -linear map has a kernel of relative size at most $1/q$, so the support induction uses $(1 - 1/q)^{k-1}$. The endpoint construction uses two independent $\mathbb{F}_q$ -linear functionals and has the same zero-product probability with q in place of p . Finally, the polarization proof only requires invertibility of $(k + 1)!$, which is ensured by $p > k + 1$, not by the weaker condition $q > k + 1$. $\square$

The character parametrization used here follows from nondegeneracy of the trace pairing: its nonzero scalar multiples give q different characters of the additive group of $\mathbb{F}_q$, hence all of them. For example, nontriviality of the trace map follows since the polynomial $X + X^p + \cdots + X^{p^{e-1}}$ is nonzero of degree less than q ; multiplication by a nonzero field element then proves nondegeneracy.

IV.24 Small characteristic and other exact boundaries

Source 36, delivered Section 10.

IV.24.1 Full symmetry does not guarantee classical integration

The symmetry theorem is valid for every odd p , regardless of k . The integration theorem is more restrictive. At the boundary $p \leq k + 1$, division by the factorial may fail for a substantive reason, not merely because the displayed formula is inconvenient.

Example IV.24.1 (Perfect energy without classical integration). On $V = \mathbb{F}_3$, let $[x] \in \{0, 1, 2\}$ be the integer representative and put

$$f(x) = \exp(2\pi i[x]/9), \quad T(a, b, c) = -abc, \quad \sigma(a, b)(x) = -abx. \quad (\text{IV.24.1})$$

Then T is symmetric, and

$$\partial_{a,b,c}f(x) = e_3(-abc), \quad \mathcal{A}_f(\sigma) = 1. \quad (\text{IV.24.2})$$

There is no function $Q : \mathbb{F}_3 \rightarrow \mathbb{F}_3$ whose third additive derivative is T .

Proof. Let $F(x) = [x]$ as an integer-valued function, computing differences modulo 9. If one increment is zero the triple difference vanishes. For three increments equal to one, write S for translation by one. Since $S^3 = I$,

$$(S - I)^3 F = 3(S - S^2)F \equiv -3 \pmod{9},$$

because $(S - S^2)F$ has values $-1, 2, -1$. For a negative increment, the difference operator obeys $\Delta_{-1} = -S^{-1}\Delta_1$. Triple differences therefore have value $-3abc$ modulo 9 for all $a, b, c \in \mathbb{F}_3$; translations of the constant triple difference do not change it. Exponentiating proves the first identity in (IV.24.2); inserting it in (IV.16.2) proves the second.

Every function $\mathbb{F}_3 \rightarrow \mathbb{F}_3$ is represented by a polynomial of degree at most two, by interpolation at the three field points. Each additive difference lowers its degree by at least one, so every third derivative is zero. Since $T(1, 1, 1) \neq 0$, no such Q exists. $\square$

The phase $[x]/9$ is a nonclassical polynomial: its third difference is constant and its fourth difference vanishes in $\mathbb{R}/\mathbb{Z}$. This example illustrates an established distinction in finite-field inverse theory, rather than discovering that distinction. Tidor's integration theorem [150] incorporates additional low-characteristic identities; it must not be read as integration of all ordinary symmetric tensors into classical phases. Also, exact integration of a prescribed derivative is stronger than merely finding correlation with some classical polynomial phase. Our example only rules out the former.

IV.24.2 Other limitations

Boundedness is a real hypothesis. Scaling f by M scales $\mathcal{A}_f$ by $|M|^{2^{k+1}}$. Without a pointwise bound, a nonsymmetric selector can therefore carry arbitrary energy. The amplitude-sensitive estimate remains correct because its right side has the same homogeneity.

One repeated function is special. The reduction to $\mathcal{Q}_B(g)$ uses the fact that all cube vertices come from derivatives of the same f . An arbitrary array of different vertex functions need not have that structure. The mixed bilinear estimate is useful for affine offsets, but it does not justify substituting our threshold into a general mixed higher-order correlation theorem.

Small energy is not exact symmetry. The endpoint construction already gives a substantial nonsymmetric energy. Below the threshold, only rank-profile consequences are proved. General approximate symmetry is a different issue: the bounded-partition-rank approximation statement proposed in Tidor’s work is false in general, as shown by Milićević’s four-linear counterexamples over $\mathbb{F}_2$ [151]. We do not relabel that disproved general statement as an open question or claim that our odd-characteristic criterion rescues it.

No end-to-end Ramsey estimate. A global selector, a finite-field domain, and sufficient energy are inputs to this manuscript. Constructing the selector from an arithmetic-progression obstruction, managing interval boundaries, and iterating a density increment are not outputs. The local result cannot by itself alter the final quantitative theorem in the repository.

IV.25 Sources 46 and 43: exact binary cubic norms and Boolean phase integration [write]

Added in the eighth write, 7 October 2026.

The next sixteen sections are source 46’s Sections 2–9 and then source 43’s Sections 2–9, on Gowers’s phase-extraction step (Lemma 17.1 and Proposition 17.2) in characteristic two, where division by the factorial is impossible and symmetry is not the whole integrability condition. *Logically the chapter stands beside source 36’s chapter* (Section IV.15), which treats odd characteristic and excludes characteristic two; it is printed here because no number printed earlier may move. Their introductions are Sections F.51 and F.48.

The common theorem, proved twice. For a symmetric trilinear form T over $\mathbb{F}_2^n$ let r be the rank of the alternating form $T(a, a, b) + T(a, b, b)$ (source 46’s C_T , source 43’s Ω_T). Then the T -twisted cubic average of eight 1-bounded functions has modulus at most $2^{-r/2}$, attained with all eight functions equal to one nonclassical cubic phase with values among the eighth roots of unity, and $r/2$ is the least codimension of a subspace on which T integrates (source 46’s Theorem IV.27.1; source 43’s Theorem IV.39.4 for one function, with an amplitude factor). The seven-function correlation is at most $2^{-r/4}$, so correlation δ forces an integrable restriction of codimension at most $2 \log_2(1/\delta)$, against $8 \log_2(1/\delta)$ in Tidor’s Proposition 4.6 (source 46’s Theorem IV.28.1; source 43’s Theorem IV.40.3). The two proofs differ: source 46 converts the cube into quadratic convolution operators and a symplectic decomposition; source 43 uses a shear identity and its division-free orthogonality theorem. Source 46’s eight-function statement contains source 43’s one-function maximum, so source 46 is printed first; source 43 is a credited second route, printed whole with its further results.

Further results. Source 46: an exact agreement threshold $(1+2^{-r/2})/2$ and a spectral stability identity (Corollary IV.28.3, Proposition IV.28.4); constructive integration and the optimal phase order; an all-characteristic scalar polarization with the optimal root-of-unity alphabet (Theorem IV.30.2); the classification of exact nonclassical integrability of the trace tensors $\text{Tr}(ch_1 \cdots h_d)$ over $\mathbb{F}_{p^n}$ with their obstruction ranks and repair costs (Theorems IV.31.1, IV.32.1, IV.32.3). Source 43: the

sharp orthogonality $\mathcal{Q}_B(g, h) \leq d_B^{-1} \|g\|_2^2 \|h\|_2^2$ on every finite abelian group, without division by two (Theorem IV.35.2); the exact symmetry threshold $1 - (1 - 1/p_*(G))^k$ on every noncyclic finite abelian group (Theorem IV.36.4), which extends source 36's threshold and its bilinear estimate; an explicit Boolean primitive (Theorem IV.37.2); the sharp certification threshold $1 - 2^{1-d}$ in every degree (Theorem IV.38.4); the fourth-power triangle bound in place of the eighth power of Tidor's Lemma 4.1 (Theorem IV.40.1); the quartic reduction and the value 11/16 on $\mathbb{F}_2^2$ (Theorems IV.41.2, IV.41.3). The four claims of source 43 that the intake did not check (orthogonality, certification threshold, fourth power, 11/16) were checked at the write (note in Section F.48). The integration criterion itself is Tao–Ziegler's and Tidor's, as both sources say; source 46 does not cite source 36, and source 43 cites it.

Relation to the formal statements (checked at HEAD 6d4942222). `lemma_17_1` and `proposition_17_2` (Sections17_18.lean, lines 38 and 48) have the exact companions `lemma_17_1_holds` (Proofs17Phase.lean, line 374) and `proposition_17_2_holds` (Proofs17PhaseRemoval.lean, line 373); they are stated on $\mathbb{Z}/N\mathbb{Z}$ with N prime and $k + 1 < N$, so characteristic two lies outside them, as both sources say. Nothing in this chapter is formalized.

IV.26 Conventions and the first integration obstruction

Source 46, delivered Section 2.

Throughout, p is prime and V is a finite-dimensional $\mathbb{F}_p$ -vector space. All averages are normalized. We write $\mathbb{T} = \mathbb{R}/\mathbb{Z}$, $e(t) = \exp(2\pi it)$, and $\iota(a) = |a|/p \pmod{1}$, where $|a| \in \{0, \dots, p-1\}$ is the standard integer representative. The map $\iota : \mathbb{F}_p \rightarrow \mathbb{T}$ is additive. A function is *1-bounded* if its modulus is at most one everywhere.

For $P : V \rightarrow \mathbb{T}$, put

$$D_h P(x) = P(x + h) - P(x).$$

A nonclassical polynomial of degree at most d is a function for which every $(d + 1)$ -fold additive derivative vanishes. A symmetric d -linear form $T : V^d \rightarrow \mathbb{F}_p$ *integrates* if

$$D_{h_1} \cdots D_{h_d} P(x) = \iota(T(h_1, \dots, h_d)) \tag{IV.26.1}$$

for all $x, h_1, \dots, h_d$. Such a P automatically has degree at most d . We refer to integration by a function with values in a coset of $p^{-R}\mathbb{Z}/\mathbb{Z}$ as integration with phase order p^R . Constants are immaterial. Classical integration means that $R = 1$ suffices.

Proposition IV.26.1 (Known integration criterion). *For symmetric d -linear T , a necessary and sufficient condition for nonclassical integration is the following repeated-variable condition whenever $d \geq p + 1$:*

$$T(\underbrace{a, \dots, a}_p, b, z_1, \dots, z_{d-p-1}) = T(a, \underbrace{b, \dots, b}_p, z_1, \dots, z_{d-p-1}). \tag{IV.26.2}$$

For $d \leq p$ there is no additional condition. A necessary condition for classical integration when $d \geq p$ is vanishing whenever p arguments coincide.

The general equivalence is the Tao–Ziegler/Tidor criterion [184, 150]; we do not reprove its general sufficiency. The necessity and the sufficiency for $d \leq p + 1$ needed throughout this article are proved below and in Section IV.29.

Lemma IV.26.2 (Repeated-variable necessity). *Every total derivative satisfies (IV.26.2).*

Proof. Let I be the augmentation ideal of the integral group ring $\mathbb{Z}[V]$, acting on functions by translations, and put $\delta_a = [a] - [0]$. Because $pa = 0$,

$$\delta_a^p = - \sum_{j=1}^{p-1} \binom{p}{j} \delta_a^j. \quad (\text{IV.26.3})$$

In particular $pI^s \subseteq I^{s+p-1} + pI^{s+1}$ for $s \geq 1$. Iterating this inclusion modulo I^{d+1} shows that pI^s annihilates every degree-at-most- d polynomial whenever $s + p - 1 > d$.

Subtract (IV.26.3) multiplied by δ_b from its version with a, b interchanged and the corresponding factor δ_a . The $j = 1$ terms cancel. Each remaining term, after multiplication by the $d - p - 1$ other differences, belongs to pI^{d-p+j} with $j \geq 2$. It therefore annihilates P . This gives $\delta_a^p \delta_b \prod \delta_{z_i} P = \delta_a \delta_b^p \prod \delta_{z_i} P$, as required. For $p = 2$ the sum of remaining terms is empty. On $\mathbb{F}_p$ -valued functions, more strongly, $\delta_a^p = [pa] - [0] = 0$, which proves the stated classical necessity. $\square$

At the first obstructed degree $d = p + 1$, define

$$C_T(a, b) = T(a, \dots, a, b) - T(a, b, \dots, b), \quad (\text{IV.26.4})$$

with p repeated arguments in each term. Symmetry and the divisibility of $\binom{p}{j}$ by p show that C_T is bilinear. It is alternating because $C_T(a, a) = 0$.

Proposition IV.26.3 (First-obstruction test). *A symmetric $(p + 1)$ -linear form integrates if and only if $C_T = 0$. For a subspace $W \leq V$, its restriction integrates if and only if $C_T|_{W \times W} = 0$.*

Proof. Necessity is Lemma IV.26.2. An explicit coordinate integrator proving sufficiency is given in Theorem IV.29.4; the same construction on a basis of W proves the restriction statement. $\square$

Lemma IV.26.4 (Alternating-form geometry). *If an alternating bilinear form C on an n -dimensional vector space has rank r , then $r = 2m$ is even. Its largest totally isotropic subspace has dimension $n - m$. Moreover, C is a sum of m elementary wedges of linear functionals, and no fewer such wedges suffice.*

Proof. If $C \neq 0$, choose u, v with $C(u, v) = 1$. Their span is nondegenerate. Subtract appropriate multiples of u, v from every other vector to split this plane orthogonally, and iterate. This gives m hyperbolic planes and a radical of dimension $n - 2m$. Taking the radical and one line from each plane gives an isotropic subspace of dimension $n - m$.

In the quotient by the radical, an isotropic subspace U is contained in its orthogonal complement, so $\dim U \leq m$. This proves maximality in V as well. The dual-coordinate expression of the hyperbolic planes gives m wedges. Each wedge has rank at most two, so rank subadditivity proves the lower bound on their number. $\square$

IV.27 The exact binary mixed cubic norm

Source 46, delivered Section 3.

In this section $V = \mathbb{F}_2^n$ and $T : V^3 \rightarrow \mathbb{F}_2$ is symmetric and trilinear. Thus

$$C_T(a, b) = T(a, a, b) + T(a, b, b), \quad r = \text{rank } C_T = 2m.$$

For eight functions $f_\varepsilon : V \rightarrow \mathbb{C}$, indexed by $\varepsilon \in \{0, 1\}^3$, define

$$\Lambda_T((f_\varepsilon)) = \mathbb{E}_{x,a,b,c} \prod_{\varepsilon \in \{0,1\}^3} \mathcal{C}^{|\varepsilon|} f_\varepsilon(x + \varepsilon_1 a + \varepsilon_2 b + \varepsilon_3 c) (-1)^{T(a,b,c)}, \quad (\text{IV.27.1})$$

where $\mathcal{C}$ denotes complex conjugation. Allowing different functions at all vertices is important: the result is stronger than a bound for one repeated function.

Theorem IV.27.1 (Sharp mixed cubic bound). *For all 1-bounded f_ε ,*

$$|\Lambda_T((f_\varepsilon))| \leq 2^{-r/2}. \quad (\text{IV.27.2})$$

The bound is attained. In fact there is a nonclassical polynomial $P : V \rightarrow \frac{1}{8}\mathbb{Z}/\mathbb{Z}$ of degree at most three such that, on setting all $f_\varepsilon = e(P)$, the average equals $2^{-r/2}$.

[write] **7 October 2026; sources 40–46.** Source 43's Theorem IV.39.4 proves the one-function case independently (with an amplitude factor), by a shear identity and its division-free orthogonality theorem; here the bound holds for eight different functions. The value $2^{-r/2}$ and the codimension $r/2$ agree.

IV.27.1 A quadratic spectral calculation

For a real-valued function κ on V , normalized convolution is

$$(K_\kappa H)(x) = \mathbb{E}_y \kappa(x + y) H(y).$$

The characters are $(-1)^{\xi(x)}$ for $\xi \in V^*$, with eigenvalues $\widehat{\kappa}(\xi) = \mathbb{E}_x \kappa(x) (-1)^{\xi(x)}$.

Lemma IV.27.2 (Flat nonzero quadratic spectrum). *Let $Q : V \rightarrow \mathbb{F}_2$ satisfy $Q(0) = 0$, with alternating polar form $C(x, h) = Q(x + h) + Q(x) + Q(h)$ of rank r . For any linear $L : V \rightarrow \mathbb{F}_2$, the function $\kappa(x) = (-1)^{Q(x) + L(x)}$ has Walsh coefficients in $\{0, 2^{-r/2}, -2^{-r/2}\}$. Consequently $\|K_\kappa\|_{2 \rightarrow 2} = 2^{-r/2}$.*

Proof. Changing variables in the square of a Walsh coefficient gives

$$\begin{aligned} |\widehat{\kappa}(\xi)|^2 &= \mathbb{E}_{x,h} (-1)^{Q(x) + Q(x+h) + (L+\xi)(h)} \\ &= \mathbb{E}_h (-1)^{Q(h) + (L+\xi)(h)} \mathbb{E}_x (-1)^{C(x,h)}. \end{aligned}$$

The inner average is the indicator of $R = \text{rad } C$. The restriction $Q|_R$ is linear. Averaging its character against $L + \xi$ therefore gives either zero or $|R|/|V| = 2^{-r}$. The coefficients are real, proving the assertion. At least one is nonzero, by Parseval. Diagonalization in the character basis proves the operator-norm formula. $\square$

IV.27.2 Proof of the upper bound

Make the invertible change of variables

$$b = a + A, \quad c = a + B, \quad y = x + a.$$

The four even-parity vertices become $x, x + A, x + B, x + A + B$; the four odd-parity vertices become $y, y + A, y + B, y + A + B$. Define

$$\begin{aligned} F_{A,B}(x) &= f_{000}(x) f_{110}(x + A) f_{101}(x + B) f_{011}(x + A + B), \\ G_{A,B}(y) &= f_{100}(y) f_{010}(y + A) f_{001}(y + B) f_{111}(y + A + B). \end{aligned}$$

Both functions have L^2 norm at most one. Multilinearity gives

$$T(a, a + A, a + B) = Q(a) + L_{A,B}(a), \quad (\text{IV.27.3})$$

where

$$Q(a) = T(a, a, a), \quad L_{A,B}(a) = T(a, a, A) + T(a, a, B) + T(a, A, B).$$

Each term in $L_{A,B}$ is linear in a : the cross terms in two equal slots cancel in characteristic two. The polar form of Q is precisely C_T . Therefore, with $\kappa_{A,B}(a) = (-1)^{Q(a)+L_{A,B}(a)}$,

$$\Lambda_T = \mathbb{E}_{A,B} \mathbb{E}_{x,y} F_{A,B}(x) \overline{G_{A,B}(y)} \kappa_{A,B}(x+y).$$

Lemma IV.27.2 and the Hilbert-space Cauchy–Schwarz inequality bound the absolute value of each inner pairing by $2^{-r/2} \|F_{A,B}\|_2 \|G_{A,B}\|_2 \leq 2^{-r/2}$. Averaging proves (IV.27.2).

IV.27.3 Construction of extremizers

For linear forms $\ell, m : V \rightarrow \mathbb{F}_2$, set

$$S_{\ell,m}(a, b, c) = \ell(a)m(b)m(c) + m(a)\ell(b)m(c) + m(a)m(b)\ell(c). \quad (\text{IV.27.4})$$

This is symmetric and

$$C_{S_{\ell,m}}(a, b) = \ell(b)m(a) + \ell(a)m(b).$$

By Lemma IV.26.4, choose $m = r/2$ independent pairs (ℓ_i, m_i) whose wedges sum to C_T . Put $S = \sum_i S_{\ell_i, m_i}$. Then $R = T - S$ has $C_R = 0$ and therefore integrates. The explicit binary construction in Section IV.29 gives an integrator $P : V \rightarrow \frac{1}{8}\mathbb{Z}/\mathbb{Z}$ of degree at most three.

For $f = e(P)$, the cubical product in (IV.27.1) equals $(-1)^{R(a,b,c)}$. Thus

$$\Lambda_T(f, \dots, f) = \mathbb{E}_{a,b,c} (-1)^{S(a,b,c)}.$$

In each symplectic coordinate pair, averaging over c leaves the condition

$$m(a)m(b) = 0, \quad \ell(a)m(b) + m(a)\ell(b) = 0.$$

Exactly eight of the sixteen choices for $(\ell(a), m(a), \ell(b), m(b))$ satisfy these two equations: four with $m(a) = m(b) = 0$, two with $(m(a), m(b)) = (1, 0)$, and two with $(0, 1)$. The block average is consequently $1/2$. The coordinate pairs are jointly independent under uniform averaging. Their product gives 2^{-m} , proving sharpness and completing Theorem IV.27.1.

Remark IV.27.3. The proof does not estimate the rank of the original trilinear tensor. It uses the rank of its repeated-variable obstruction. Tensors of very different ordinary tensor ranks can therefore have exactly the same twisted cube norm.

IV.28 Restriction, exact agreement, and stability

Source 46, delivered Section 4.

IV.28.1 A factor-four improvement of the symmetric seven-function bound

For seven 1-bounded functions indexed by the nonempty subsets of $\{1, 2, 3\}$, put

$$\Gamma_T = \mathbb{E}_{a,b,c} b_1(a)b_2(b)b_3(c)b_{12}(a+b)b_{13}(a+c) \cdot b_{23}(b+c)b_{123}(a+b+c)(-1)^{T(a,b,c)}. \quad (\text{IV.28.1})$$

Theorem IV.28.1 (Improved symmetric restriction estimate). *Let $T : V^3 \rightarrow \mathbb{F}_2$ be symmetric and trilinear, with $\text{rank } C_T = r$. Then*

$$|\Gamma_T| \leq 2^{-r/4}. \quad (\text{IV.28.2})$$

If $|\Gamma_T| \geq \delta > 0$, there is a subspace $W \leq V$ with

$$\text{codim } W = \frac{r}{2} \leq 2 \log_2(1/\delta) \quad (\text{IV.28.3})$$

on which T integrates to a nonclassical cubic polynomial. Among all integrable restrictions of this particular T , the codimension $r/2$ is minimal.

Proof. Cauchy–Schwarz in (b, c) removes the product $b_2(b)b_3(c)b_{23}(b+c)$ and gives

$$|\Gamma_T|^2 \leq \mathbb{E}_{b,c} \left| \mathbb{E}_a b_1(a)b_{12}(a+b)b_{13}(a+c)b_{123}(a+b+c)(-1)^{T(a,b,c)} \right|^2.$$

Expand the square with the second variable $a+h$. The tensor factor becomes $(-1)^{T(h,b,c)}$. The resulting expression is an eight-function cube average with basepoint a , directions (h, b, c) , and assignments

$$\begin{aligned} f_{000} &= f_{100} = b_1, & f_{010} &= f_{110} = \overline{b_{12}}, \\ f_{001} &= f_{101} = \overline{b_{13}}, & f_{011} &= f_{111} = b_{123}. \end{aligned}$$

It is bounded by $2^{-r/2}$ by Theorem IV.27.1. Taking square roots proves (IV.28.2). Hence $r \leq 4 \log_2(1/\delta)$. A largest isotropic subspace for C_T has codimension $r/2$ by Lemma IV.26.4, and is exactly an integrable restriction by Proposition IV.26.3. These same facts give minimality. $\square$

[write] **7 October 2026; sources 40–46.** Source 43’s Theorem IV.40.3 gives the same bound $2^{-r/4}$ and codimension $\lceil 2 \log_2(1/\delta) \rceil$ through its fourth-power triangle estimate. Neither source claims the coefficient two optimal.

Comparison with the published bound. Tidor’s Proposition 4.6 [150] has the identical seven-function hypothesis and concludes $\text{codim } W \leq 8 \log_2(1/\delta)$. There are two improvements in the preceding argument: the spectral cube estimate gives $r \leq 4 \log_2(1/\delta)$, and one restricts to a maximal isotropic subspace rather than to the full radical. The latter saves another factor of two in codimension. No symmetry assumption has been dropped. For a nonsymmetric tensor one must still pay for symmetrization; we do not claim here a replacement for the constants in Tidor’s nonsymmetric Corollary 4.8 or in a full inverse theorem. Nor does sharpness of Theorem IV.27.1 establish sharpness of (IV.28.2): one Cauchy–Schwarz step intervenes.

IV.28.2 Exact derivative-spectrum maximum

For a single 1-bounded $f : V \rightarrow \mathbb{C}$, define

$$M_T(f) = \mathbb{E}_{a,b} \left| \mathbb{E}_x f(x) \overline{f(x+a)} \overline{f(x+b)} f(x+a+b) (-1)^{T(a,b,x)} \right|^2. \quad (\text{IV.28.4})$$

This is a squared Fourier-coefficient moment of the second multiplicative derivative, evaluated at the linear frequency $x \mapsto T(a, b, x)$.

Corollary IV.28.2 (Exact cubic spectral moment). *For $r = \text{rank } C_T$,*

$$\sup_{\|f\|_\infty \leq 1} M_T(f) = 2^{-r/2}.$$

The supremum is attained by the eighth-root phase in Theorem IV.27.1.

Proof. In (IV.28.4), expand the square with the second basepoint $x + c$. The product is precisely the eight-vertex product in (IV.27.1), and $T(a, b, x) + T(a, b, x + c) = T(a, b, c)$. Thus $M_T(f) = \Lambda_T(f, \dots, f)$. Apply Theorem IV.27.1. $\square$

IV.28.3 Optimal approximate polarization

Let $P_\varepsilon : V \rightarrow \mathbb{T}$ be completely arbitrary functions, not assumed to be polynomials. For uniform (x, a, b, c) , compare

$$\sum_{\varepsilon \in \{0,1\}^3} (-1)^{|\varepsilon|} P_\varepsilon(x + \varepsilon_1 a + \varepsilon_2 b + \varepsilon_3 c) \quad \text{with} \quad \iota(T(a, b, c)). \quad (\text{IV.28.5})$$

Equality is in $\mathbb{T}$.

Corollary IV.28.3 (Exact agreement and squared-error thresholds). *Writing $m = \text{rank } C_T/2$, the maximum probability of equality in (IV.28.5), over all eight functions P_ε , is*

$$\frac{1 + 2^{-m}}{2}. \quad (\text{IV.28.6})$$

The minimum mean squared distance between the corresponding unit complex numbers is

$$2(1 - 2^{-m}). \quad (\text{IV.28.7})$$

Both optima are attained with all P_ε equal to one nonclassical cubic with values in $\frac{1}{8}\mathbb{Z}/\mathbb{Z}$.

Proof. Let Z be the exponential of the difference between the two sides of (IV.28.5). Then $|Z| = 1$ and $|\mathbb{E}Z| \leq 2^{-m}$ by Theorem IV.27.1. If $s = \Pr(Z = 1)$, then $\text{Re } \mathbb{E}Z \geq 2s - 1$. This proves the upper bound for agreement. Also $\mathbb{E}|Z - 1|^2 = 2 - 2\text{Re } \mathbb{E}Z \geq 2(1 - 2^{-m})$. For the extremizer constructed above, $Z = (-1)^{S(a,b,c)}$ and $\mathbb{E}Z = 2^{-m}$. Since Z takes only the values $1, -1$, equality holds in both bounds. $\square$

Thus independent choices at the vertices cannot conceal a genuine integration obstruction. The conclusion quantifies this failure, whereas a formal factorial calculation only indicates that one particular formula is unavailable.

IV.28.4 A spectral stability identity

The proof of Theorem IV.27.1 retains information about near equality. For a common function f , let

$$F_{A,B}(x) = f(x)f(x+A)f(x+B)f(x+A+B), \quad \lambda = 2^{-r/2}.$$

The kernel $\kappa_{A,B}$ in (IV.27.3) has three possible eigenvalues: $+\lambda, 0, -\lambda$. Let $\Pi_{A,B}^+, \Pi_{A,B}^0, \Pi_{A,B}^-$ be the orthogonal projections onto these three Walsh spectral subspaces. Set

$$A_f = \mathbb{E}_{A,B} \|F_{A,B}\|_2^2, \quad \mu_\sigma = \mathbb{E}_{A,B} \left\| \Pi_{A,B}^\sigma F_{A,B} \right\|_2^2 \quad (\sigma \in \{+, 0, -\}).$$

Proposition IV.28.4 (Exact defect decomposition). *For every 1-bounded f ,*

$$1 - \frac{M_T(f)}{\lambda} = (1 - A_f) + \mu_0 + 2\mu_-. \quad (\text{IV.28.8})$$

In particular, if $M_T(f) \geq (1 - \varepsilon)\lambda$, then

$$\begin{aligned} \mathbb{E}_x |f(x)|^2 &\geq 1 - \varepsilon, \\ \mathbb{E}_{A,B} \left\| F_{A,B} - \Pi_{A,B}^+ F_{A,B} \right\|_2^2 &\leq \varepsilon, \\ \mu_- &\leq \varepsilon/2. \end{aligned}$$

Proof. For a common function, the two functions $F_{A,B}, G_{A,B}$ in the upper-bound proof coincide. Parseval and the real eigenvalues give $M_T(f) = \lambda(\mu_+ - \mu_-)$, whereas $A_f = \mu_+ + \mu_0 + \mu_-$. Subtraction proves (IV.28.8). Every term on its right is nonnegative because $A_f \leq 1$. Moreover $A_f \leq \mathbb{E}|f|^2$, since three of the four factors in $|F_{A,B}(x)|^2$ can be bounded by one. The conclusions follow. $\square$

This is stability of the transformed four-point functions in a completely specified spectral space. It is not yet a classification of f itself near a polynomial-phase extremizer; recovering such a classification is a separate problem.

IV.29 Constructive integration and optimal phase order

Source 46, delivered Section 5.

We now supply the explicit integrators used above, and formulate an all-characteristic scalar replacement for Gowers' polarization identity. This also distinguishes higher-root phenomena from true obstructions.

IV.29.1 A finite-difference calculation on the prime field

For $1 \leq a \leq p-1$, write $L_a(x) = |x|^a$ on $\mathbb{F}_p$, viewed first as an integer-valued periodic function. Let $\delta f(x) = f(x+1) - f(x)$.

Lemma IV.29.1 (Valuation and leading difference). *For $k \geq 0$,*

$$v_p(\delta^k L_a(x)) \geq \max \left\{ 0, \left\lceil \frac{k-a}{p-1} \right\rceil \right\} \quad (\text{IV.29.1})$$

uniformly in x , with $v_p(0) = +\infty$. If $d = a + s(p-1)$, then

$$\delta^d L_a(x) \equiv (-1)^s p^s a! \pmod{p^{s+1}}, \quad (\text{IV.29.2})$$

$$\delta^{d+1} L_a(x) \equiv 0 \pmod{p^{s+1}}. \quad (\text{IV.29.3})$$

Proof. For $k < p$, ordinary polynomial differencing modulo p proves (IV.29.1): when $k > a$, the reduction modulo p vanishes. Since the translation has order p , $(1 + \delta)^p = 1$, hence

$$\delta^k = - \sum_{j=1}^{p-1} \binom{p}{j} \delta^{k-p+j} \quad (k \geq p).$$

Strong induction on k proves (IV.29.1), using $p \mid \binom{p}{j}$ and observing that the smallest resulting valuation is the $j = 1$ bound. For (IV.29.2), use induction on s . At $s = 0$ the familiar a th difference of x^a is $a!$ modulo p . At $d = a + s(p-1)$, the $j = 1$ term is $-p\delta^{d-(p-1)}L_a$. Every term with $j \geq 2$ is divisible by p^{s+1} by (IV.29.1), giving the induction step. Equation (IV.29.3) is another direct instance of that valuation bound. $\square$

Theorem IV.29.2 (Optimal scalar integrator in every degree). *Given $d \geq 1$, write uniquely*

$$d = a + s(p-1), \quad 1 \leq a \leq p-1, \quad s \geq 0.$$

Choose an integer u with $u(-1)^s a! \equiv 1 \pmod{p}$. Then

$$Q_d(x) = \frac{u|x|^a}{p^{s+1}} \pmod{1} \quad (\text{IV.29.4})$$

has degree at most d and satisfies

$$D_{h_1} \cdots D_{h_d} Q_d(x) = \iota(h_1 \cdots h_d). \quad (\text{IV.29.5})$$

The smallest possible phase order for any integrator of this nonzero tensor is

$$p^{s+1} = p^{\lceil d/(p-1) \rceil}. \quad (\text{IV.29.6})$$

The lower bound allows arbitrary functions, provided their values lie in a coset of the indicated finite subgroup of $\mathbb{T}$.

Proof. Every difference D_h on $\mathbb{F}_p$ is an integer-polynomial multiple of δ . Lemma IV.29.1 therefore shows that Q_d has degree at most d . Its top derivative is independent of x and additive in each direction: the identity $D_{h+k} = D_h + D_k + D_h D_k$ leaves an error of order $d+1$. As an additive map out of $\mathbb{F}_p$ in each slot, it is determined by its value on $(1, \dots, 1)$, which is $1/p$ by (IV.29.2) and the choice of u . This proves (IV.29.5).

For optimality, subtract a constant from any proposed integrator with values in a coset of $p^{-R}\mathbb{Z}/\mathbb{Z}$. Its numerator is a function $\mathbb{F}_p \rightarrow \mathbb{Z}/p^R\mathbb{Z}$. The evaluation matrix for $1, |x|, \dots, |x|^{p-1}$ is a Vandermonde matrix whose determinant is a unit modulo p , and hence modulo p^R . These functions therefore interpolate all possible numerators. By (IV.29.1), each $|x|^a/p^R$ has degree at most $a + (R-1)(p-1) \leq R(p-1)$; the constant has degree zero. Thus every such function has degree at most $R(p-1)$. A nonzero d th derivative requires $d \leq R(p-1)$, proving the lower bound. $\square$

Example IV.29.3 (Binary scalar phases). For $p = 2$, one may take $Q_d(x) = (-1)^{d-1}|x|/2^d$. In particular a nonzero cubic tensor on $\mathbb{F}_2$ requires eighth roots, even though its total derivative takes only the values $0, 1/2$. The issue is phase depth, not a failure of integrability.

IV.29.2 Multivariate coordinate integration through degree $p + 1$

We will also use the following multivariate consequence of the valuation argument. If $0 \leq a_i < p$ and $A = \sum_i a_i > 0$, then

$$x \mapsto \frac{\prod_i |x_i|^{a_i}}{p^{s+1}} \quad \text{has degree at most } A + s(p-1). \quad (\text{IV.29.7})$$

Indeed the coordinate augmentation ideal is generated by the coordinate differences δ_i . A product $\prod_i \delta_i^{k_i}$ acts separately on the factors. The sum of the lower bounds $\max(0, \lceil (k_i - a_i)/(p-1) \rceil)$ is at least $s + 1$ whenever $\sum_i k_i > A + s(p-1)$. This proves (IV.29.7); coordinates with $a_i = 0$ cause no difficulty, since any difference in such a coordinate annihilates that factor.

Fix a basis $e_1, \dots, e_n$ of V . For $\alpha \in \mathbb{Z}_{\geq 0}^n$ with $|\alpha| = d$, let

$$b_\alpha = T(e_1^{\alpha_1}, \dots, e_n^{\alpha_n}),$$

where superscripts denote repetitions, not multiplication. A symmetric multilinear form is determined by these coefficients. For coefficients in formulas below, division by a factorial means inversion in $\mathbb{F}_p$, followed by the standard integer representative.

Theorem IV.29.4 (Explicit low-degree integration). *Every symmetric d -linear form with $d \leq p$ integrates. A symmetric $(p+1)$ -linear form integrates if $C_T = 0$. Integrators can be constructed as follows.*

For every α with all $\alpha_i < p$, include the classical term

$$\frac{1}{p} \left| \frac{b_\alpha}{\prod_i \alpha_i!} \right| \prod_i |x_i|^{\alpha_i}. \quad (\text{IV.29.8})$$

When $d = p$, include additionally $|b_{pe_i}| Q_p(x_i)$ for each pure coefficient. When $d = p + 1$, instead include the pure terms $|b_{(p+1)e_i}| Q_{p+1}(x_i)$ and, for every $i < j$, the term

$$- \frac{|b_{pe_i+e_j}| |x_i| |x_j|}{p^2}. \quad (\text{IV.29.9})$$

The last prescription is valid because $C_T = 0$ makes $b_{pe_i+e_j} = b_{e_i+pe_j}$. Summing these terms modulo one gives an integrator.

Proof. The total derivative of (IV.29.8) has exactly the prescribed α component, by ordinary differentiation in the prime-field polynomial algebra; all the relevant factorials are invertible. The scalar terms are handled by Theorem IV.29.2. Their derivatives vanish whenever a direction from another coordinate occurs.

It remains to check the paired term. Modulo p^2 , the recurrence for δ_i^p acting on $|x_i|$ has only the term $-p\delta_i|x_i|$ remaining: for $2 \leq j < p$, $\delta_i^j|x_i|$ is divisible by p . Thus the $pe_i + e_j$ and $e_i + pe_j$ components of $-|x_i||x_j|/p^2$ are both $1/p$. Its pure components vanish by the scalar degree bound. If both multiplicities are between 2 and $p-1$, each coordinate difference contributes a factor p , so that component vanishes modulo one. Any component involving a third coordinate vanishes identically. The degree is at most $p+1$ by (IV.29.7).

For $d = p+1$ these are all possible multiplicity patterns: a pure $(p+1)$, a pair $(p,1)$, or multiplicities all smaller than p . The equality of paired coefficients is exactly $C_T(e_i, e_j) = 0$. Matching every basis component proves the theorem. $\square$

Corollary IV.29.5 (The binary cubic integrator). *If $R : V^3 \rightarrow \mathbb{F}_2$ is symmetric and $C_R = 0$, then, writing $r_i = R(e_i, e_i, e_i)$, $r_{ij} = R(e_i, e_i, e_j) = R(e_i, e_j, e_j)$, and $r_{ijk} = R(e_i, e_j, e_k)$,*

$$P(x) = \sum_i \frac{r_i |x_i|}{8} - \sum_{i < j} \frac{r_{ij} |x_i| |x_j|}{4} + \sum_{i < j < k} \frac{r_{ijk} |x_i| |x_j| |x_k|}{2} \pmod{1} \quad (\text{IV.29.10})$$

has $D_a D_b D_c P = \iota(R(a, b, c))$.

The coefficients are binary representatives. This formula makes the extremizer construction in Section IV.27 completely effective.

IV.30 An all-characteristic replacement of scalar polarization

Source 46, delivered Section 6.

IV.30.1 Independent vertex functions do not remove exact obstructions

The following observation applies in every characteristic and does not require any regularity of the vertex functions.

Lemma IV.30.1 (Vertex separation forces integration). *Let $T : V^d \rightarrow \mathbb{F}_p$ be multilinear. Suppose functions $P_\varepsilon : V \rightarrow \mathbb{T}$ satisfy*

$$\iota(T(h_1, \dots, h_d)) = \sum_{\varepsilon \in \{0,1\}^d} (-1)^{|\varepsilon|} P_\varepsilon(s - \varepsilon \cdot h) \quad (\text{IV.30.1})$$

for all $s, h_1, \dots, h_d$. Then $P_{(1,\dots,1)}$ integrates T . Conversely, an integrator gives (IV.30.1) by taking every vertex function equal to it.

Proof. For each i , apply the operator that subtracts the expression with $h_i = 0$. The left side is unchanged by multilinearity. Every vertex missing at least one direction is annihilated. For the all-ones vertex, the alternating expansion that remains is

$$\sum_{J \subseteq [d]} (-1)^{|J|} P_{(1,\dots,1)} \left(s - \sum_{j \in J} h_j \right) = D_{h_1} \cdots D_{h_d} P_{(1,\dots,1)} \left(s - \sum_i h_i \right).$$

This proves the first assertion. The same identity proves the converse. □

IV.30.2 Scalar multi-affine polarization with optimal roots

Let

$$\sigma(x_1, \dots, x_k) = \sum_{J \subseteq [k]} c_J \prod_{j \in J} x_j \quad (c_J \in \mathbb{F}_p).$$

This is separately affine in every variable; the constant monomial is allowed. Put $Q_0 = 1/p$, and use the functions Q_j of Theorem IV.29.2 for $j \geq 1$.

Theorem IV.30.2 (All-characteristic scalar polarization). *Define, for $\varepsilon \in \{0, 1\}^k$,*

$$P_\varepsilon(s) = \sum_{J \supseteq \text{supp } \varepsilon} |c_J| Q_{|J|}(s). \quad (\text{IV.30.2})$$

Then

$$\iota(\sigma(x)) = \sum_{\varepsilon \in \{0,1\}^k} (-1)^{|\varepsilon|} P_\varepsilon(s - \varepsilon \cdot x) \quad (\text{IV.30.3})$$

for all x, s . For $k \geq 1$, all P_ε have degree at most k and values in $p^{-R}\mathbb{Z}/\mathbb{Z}$, where $R = \lceil k/(p-1) \rceil$. This exponent is the optimal universal choice. If $k < p$, the construction can be taken entirely in $p^{-1}\mathbb{Z}/\mathbb{Z}$.

Proof. For a fixed J , its contribution to the right side is

$$|c_J| \sum_{\text{supp } \varepsilon \subseteq J} (-1)^{|\varepsilon|} Q_{|J|}(s - \varepsilon \cdot x) = \iota \left(c_J \prod_{j \in J} x_j \right)$$

by (IV.29.5), including the empty-set case. Summing proves the identity. The degree and phase-order claims follow from Theorem IV.29.2. For the monomial $\prod_{i=1}^k x_i$, Lemma IV.30.1 makes the all-ones vertex an integrator. The lower bound in Theorem IV.29.2 therefore proves optimality even among unrelated vertex functions with a common finite phase alphabet. $\square$

This removes the condition $k < p$ from the *existence of a phase-valued* scalar polarization identity. It does not remove it while retaining the original classical polynomial codomain. Conflating those statements would produce an incorrect extension of Gowers' lemma.

IV.31 Complete classification of trace-multiplication tensors

Source 46, delivered Section 7.

Let $F = \mathbb{F}_q$ with $q = p^n$, viewed additively as an n -dimensional $\mathbb{F}_p$ -space. For $c \in F^\times$ and $d \geq 1$, set

$$T_{c,d}(h_1, \dots, h_d) = \text{Tr}_{F/\mathbb{F}_p}(ch_1 \cdots h_d). \quad (\text{IV.31.1})$$

It is symmetric and $\mathbb{F}_p$ -multilinear. The trace pairing is nondegenerate: $\text{Tr}(uv) = 0$ for every $v \in F$ implies $u = 0$. We use also the Frobenius adjoint identity

$$\text{Tr}(ub^p) = \text{Tr}(u^{p^{n-1}}b), \quad (\text{IV.31.2})$$

which follows from invariance of trace under Frobenius.

Theorem IV.31.1 (Exact trace integration classification). *For nonzero c , the tensor $T_{c,d}$ integrates precisely in the following cases:*

<i>Extension degree</i>	<i>Tensor degree</i>	<i>Integrability condition</i>
$n = 1$	<i>every $d \geq 1$</i>	<i>always</i>
$n > 1$	$d \leq p$	<i>always</i>
$n > 1$	$d = p + 1$	$n = 2$ and $c \in \mathbb{F}_p^\times$
$n > 1$	$d \geq p + 2$	<i>never</i>

For all n , classical integration is possible exactly when $d < p$.

Proof. For $n = 1$, Theorem IV.29.2 applies, multiplied by the nonzero coefficient. For $d \leq p$, use Theorem IV.29.4.

At $d = p + 1$, the obstruction is

$$C_c(a, b) = \text{Tr}(c(a^p b - ab^p)). \quad (\text{IV.31.3})$$

By (IV.31.2) and nondegeneracy, $C_c = 0$ is equivalent to

$$ca^p = (ca)^{p^{n-1}} \quad \text{for every } a \in F.$$

Raising to the p th power gives the equivalent condition

$$c^p a^{p^2} = ca \quad \text{for every } a \in F. \quad (\text{IV.31.4})$$

Taking $a = 1$ forces $c^p = c$. With $c \neq 0$, the remaining condition is $a^{p^2} = a$ for all a , which is equivalent to $n \mid 2$. For $n > 1$ this gives exactly $n = 2$ and $c \in \mathbb{F}_p^\times$. Sufficiency follows from Proposition IV.26.3.

If $d \geq p + 2$, apply the necessary repeated-variable identity and set all but one of the extra directions equal to one. Their product z is arbitrary. The identity requires $C_{cz} = 0$ for every $z \in F$. For $n > 1$, the preceding classification shows that not every nonzero coefficient has vanishing obstruction. Since $cF^\times = F^\times$, this is impossible.

For classical integration with $d \geq p$, a total derivative must vanish when p directions are equal. Setting all remaining directions to one would force $\text{Tr}(ca^p) = 0$ for every a . Frobenius is bijective and the trace pairing is nondegenerate, a contradiction. Conversely, for $d < p$ the classical terms (IV.29.8) suffice. $\square$

Corollary IV.31.2 (Optimal phase orders in the admissible trace cases). *For $n = 1$ the minimum phase order is $p^{\lceil d/(p-1) \rceil}$. For $n > 1$, it is p when $d < p$, and p^2 when $d = p$. In the exceptional case $n = 2$, $d = p + 1$, $c \in \mathbb{F}_p^\times$, it is also p^2 , including $p = 2$.*

Proof. The scalar assertion is Theorem IV.29.2. For $n > 1$ and $d \leq p$, the coordinate construction gives the upper bounds, and the classical nonintegrability in Theorem IV.31.1 proves optimality at $d = p$.

For the exceptional case with $p \geq 3$, the pure term Q_{p+1} already has phase order p^2 , so the full coordinate construction does also. If $p = 2$, the potentially eighth-root pure terms have coefficient $\text{Tr}(ca^3) = 0$ for every $a \in \mathbb{F}_4$, because a^3 is the norm into $\mathbb{F}_2$ and the trace of an element of $\mathbb{F}_2$ from $\mathbb{F}_4$ is zero. Only the paired fourth-root terms remain. The impossibility of classical integration proves the lower bound in both cases. $\square$

Example IV.31.3 (The smallest extension distinguishes three phenomena). Let $F = \mathbb{F}_4 = \mathbb{F}_2[\theta]/(\theta^2 + \theta + 1)$, and write $z = x + y\theta$. For $c = 1$, the cubic trace tensor integrates to

$$P(z) = -|x||y|/4 \pmod{1}.$$

Thus fourth roots suffice, even though a nonzero scalar cubic on $\mathbb{F}_2$ requires eighth roots. For $c = \theta$, however,

$$C_\theta(1, \theta) = \text{Tr}(\theta^2) - \text{Tr}(\theta^3) = 1.$$

No circle-valued cubic integrator exists, regardless of phase alphabet. The mixed cube norm is nevertheless exactly $1/2$, by Theorem IV.27.1. An explicit common-function extremizer is

$$f(x + y\theta) = \exp\left(\frac{2\pi i(|x| + |y|)}{8}\right).$$

Its cubic phase integrates the residual diagonal tensor, not T_θ itself; the remaining elementary obstruction block has bias $1/2$. This separates classical failure, higher-root repair, and genuine nonintegrability in one very small field.

IV.32 Exact restriction dimensions and repair costs

Source 46, delivered Section 8.

IV.32.1 The obstruction rank for trace tensors

Continue at the first obstructed degree $d = p + 1$, and write C_c for (IV.31.3).

Theorem IV.32.1 (Explicit radical and maximal integrable subspace). *Let $t_c = \dim_{\mathbb{F}_p} \text{rad } C_c$. Then*

$$t_c = \begin{cases} 1, & n \text{ odd}, \\ 2, & n \text{ even and } c^{(q-1)/(p+1)} = 1, \\ 0, & n \text{ even and } c^{(q-1)/(p+1)} \neq 1. \end{cases} \quad (\text{IV.32.1})$$

Consequently

$$\text{rank } C_c = n - t_c, \quad \max_{W: T_{c,p+1}|_W \text{ integrates}} \dim W = \frac{n + t_c}{2}. \quad (\text{IV.32.2})$$

For even n , precisely $(q-1)/(p+1)$ nonzero coefficients have $t_c = 2$; the other $p(q-1)/(p+1)$ have $t_c = 0$.

Proof. The trace-pairing calculation leading to (IV.31.4), now with a fixed, shows that the radical is the kernel of

$$a \longmapsto c^p a^{p^2} - ca.$$

A nonzero radical element therefore satisfies

$$a^{p^2-1} = c^{1-p}. \quad (\text{IV.32.3})$$

In the cyclic group $F^\times$, the power map on the left has kernel of size $\gcd(p^2 - 1, p^n - 1) = p^{\gcd(2,n)} - 1$.

If n is odd, this size is $p - 1$ and the solvability condition is automatic: $(c^{1-p})^{(q-1)/(p-1)} = 1$. There are $p - 1$ nonzero solutions and the zero solution, so the radical has dimension one. If n is even, the kernel size is $p^2 - 1$, and solvability is equivalent to $c^{(q-1)/(p+1)} = 1$. There are then p^2 radical elements including zero; otherwise there is only zero. This proves (IV.32.1).

The rank formula follows. By Proposition IV.26.3, an integrable restriction is precisely an isotropic subspace for C_c . Lemma IV.26.4 proves the maximum in (IV.32.2). Finally the coefficient condition defines a subgroup of $F^\times$ of order $(q-1)/(p+1)$, giving the counts. $\square$

Corollary IV.32.2 (Explicit sharp binary trace moments). *For $p = 2$, $q = 2^n$, and $T_c(a, b, x) = \text{Tr}(cabx)$, the moment (IV.28.4) satisfies*

$$\sup_{\|f\|_\infty \leq 1} M_{T_c}(f) = \begin{cases} 2^{-(n-1)/2}, & n \text{ odd}, \\ 2^{-(n-2)/2}, & n \text{ even}, c^{(2^n-1)/3} = 1, \\ 2^{-n/2}, & n \text{ even}, c^{(2^n-1)/3} \neq 1. \end{cases} \quad (\text{IV.32.4})$$

Every value in this formula is attained by an eighth-root cubic phase.

Proof. Combine Corollary IV.28.2 and Theorem IV.32.1. □

n	Number of coefficients	r	Largest $\dim W$	$\sup M_{T_c}$
2	1	0	2	1
2	2	2	1	1/2
3	7	2	2	1/2
4	5	2	3	1/2
4	10	4	2	1/4
5	31	4	3	1/4
6	21	4	4	1/4
6	42	6	3	1/8
7	127	6	4	1/8
8	85	6	5	1/8
8	170	8	4	1/16

The table is a specialization of the proved formulas, not a numerical extrapolation. It shows that exact integration of arbitrary trace tensors cannot be recovered by a dimension-independent codimension loss. This does not contradict inverse theorems: the associated maximal correlations become small as the required codimension increases.

IV.32.2 An exact elementary symmetric repair invariant

For any prime p , and $\ell, m \in V^*$, define the symmetric $(p+1)$ -linear form

$$S_{\ell, m}(h_1, \dots, h_{p+1}) = \sum_{j=1}^{p+1} \ell(h_j) \prod_{t \neq j} m(h_t). \quad (\text{IV.32.5})$$

A direct evaluation, using $u^p = u$ in $\mathbb{F}_p$, gives

$$C_{S_{\ell, m}}(a, b) = \ell(b)m(a) - \ell(a)m(b). \quad (\text{IV.32.6})$$

Define the *elementary symmetric repair number* of a symmetric $(p+1)$ -linear T as the least number of forms (IV.32.5) whose subtraction makes T integrable. Scalar coefficients are absorbed into ℓ .

Theorem IV.32.3 (Exact elementary repair number). *For every symmetric $(p+1)$ -linear T ,*

$$\text{elementary symmetric repair number}(T) = \frac{\text{rank } C_T}{2} = \min\{\text{codim } W : T|_W \text{ integrates}\}. \quad (\text{IV.32.7})$$

In particular, for $T_{c, p+1}$ this number is $(n - t_c)/2$ with t_c given by (IV.32.1).

Proof. Write C_T as a sum of $r/2$ wedges by Lemma IV.26.4, choosing their orientations to match (IV.32.6). Subtract the corresponding $S_{\ell,m}$. The residual form has zero obstruction and integrates by Theorem IV.29.4. Conversely, each elementary form has obstruction rank at most two. Subtracting fewer than $r/2$ cannot cancel a rank- r obstruction, by rank subadditivity. The equality with minimum codimension follows again from the isotropic characterization. $\square$

IV.32.3 Ordinary slice-rank repair is related but different

A slice of a d -linear form is a term that is linear in one chosen slot and multilinear in the other $d - 1$ slots. The slice rank is the minimum number of such terms. Define

$$\rho(T) = \min\{\text{srank}(T - R) : R \text{ is symmetric and integrates}\}.$$

Proposition IV.32.4 (Slice-rank repair bounds). *For symmetric $(p + 1)$ -linear T with $r = \text{rank } C_T$,*

$$r/2 \leq \rho(T) \leq r. \quad (\text{IV.32.8})$$

If $r = 2$, then $\rho(T) = 2$. For $p = 2$, slice rank here equals partition rank.

Proof. Each elementary correction in (IV.32.5) has slice rank at most two: separate the $j = 1$ term, and factor $m(h_1)$ from all remaining terms. Theorem IV.32.3 gives the upper bound.

Suppose $T - R$ is a sum of s slices. Intersect the kernels of all s singleton linear factors, and call this subspace U . Then $\text{codim } U \leq s$ and $(T - R)|_U = 0$. Since $R|_U$ integrates, $T|_U$ does also. The exact minimum-codimension statement gives $s \geq r/2$.

To handle $r = 2$, a nonzero symmetric slice-rank-one tensor must be a scalar multiple of a pure power $\ell^{\otimes d}$. Indeed, write it as $\ell(h_1)B(h_2, \dots, h_d)$ and choose v with $\ell(v) = 1$. Symmetry in the first two slots shows that B is divisible by ℓ in its first slot; repeating this argument proves the pure-power assertion. For $d = p + 1$, such a tensor has zero obstruction, because $\ell(a)^p \ell(b) = \ell(a)\ell(b)^p$. Thus a rank-two obstruction cannot be repaired by one slice. The upper bound is two. Finally, every bipartition of three slots has a singleton side, so slice and partition rank coincide for trilinear tensors. $\square$

The factor-two interval in (IV.32.8) is not asserted to close in general. The exactly determined quantity in Theorem IV.32.3 is the *specified elementary symmetric* repair number, not unrestricted tensor rank or unrestricted partition rank in higher degree.

IV.33 Returning to Gowers' derivative-spectrum step

Source 46, delivered Section 9.

The preceding results concern exactly the type of polynomial untwisting that appears near the end of Gowers' argument. We spell out the analytic translation to avoid confusing an integration theorem with a full inverse theorem.

IV.33.1 A homogeneous frequency gives an exact norm identity

For $h = (h_1, \dots, h_j)$, use the backward-cube convention

$$\mathcal{B}_h f(x) = \prod_{\varepsilon \in \{0,1\}^j} \mathcal{C}^{|\varepsilon|} f(x - \varepsilon \cdot h).$$

This convention makes the additive phase identity

$$\sum_{\varepsilon} (-1)^{|\varepsilon|} P(x - \varepsilon \cdot h) = D_{h_1} \cdots D_{h_j} P \left(x - \sum_i h_i \right)$$

valid with a positive sign. It gives the usual Gowers norm after averaging: $\|f\|_{U^j}^{2^j} = \mathbb{E}_{x,h} \mathcal{B}_h f(x)$.

Let $F = \mathbb{F}_{p^n}$, $\chi(z) = e(\text{Tr}(z)/p)$, and $d \geq 2$. Define

$$\mathcal{M}_{c,d}(f) = \mathbb{E}_{h_1, \dots, h_{d-1}} \left| \mathbb{E}_x \mathcal{B}_{(h_1, \dots, h_{d-1})} f(x) \overline{\chi(ch_1 \cdots h_{d-1}x)} \right|^2. \quad (\text{IV.33.1})$$

Proposition IV.33.1 (Exact homogeneous untwisting). *If P integrates $T_{c,d}$, then for every function $f : F \rightarrow \mathbb{C}$,*

$$\mathcal{M}_{c,d}(f) = \|fe(-P)\|_{U^d}^{2^d}. \quad (\text{IV.33.2})$$

Proof. Expand the square in (IV.33.1) and write its second basepoint as $x - h_d$. The character factor is $\overline{\chi(ch_1 \cdots h_d)}$, and the function product is $\mathcal{B}_{(h_1, \dots, h_d)} f(x)$. Therefore

$$\mathcal{M}_{c,d}(f) = \mathbb{E}_{x, h_1, \dots, h_d} \mathcal{B}_h f(x) e(-\iota(T_{c,d}(h))).$$

The defining equation for P identifies the integrand with $\mathcal{B}_h(fe(-P))(x)$. Averaging proves (IV.33.2). $\square$

In admissible trace cases, this is an equality rather than merely a correlation-to-norm implication. In nonintegrable cases there is no exact vertex-separated replacement, by Lemma IV.30.1. For binary cubics the substitute is not an unspecified error: it is the exact ceiling in Corollary IV.32.2.

IV.33.2 General scalar multi-affine frequencies

For a separately affine $\sigma : \mathbb{F}_p^{d-1} \rightarrow \mathbb{F}_p$, define

$$\mathcal{M}_{\sigma}(f) = \mathbb{E}_{h \in \mathbb{F}_p^{d-1}} \left| \mathbb{E}_x \mathcal{B}_h f(x) e(-\iota(\sigma(h)x)) \right|^2.$$

The phase after squaring is $\iota(h_d \sigma(h_1, \dots, h_{d-1}))$, which is separately affine in the d directions.

Theorem IV.33.2 (Scalar phase-valued extension of the untwisting step). *For $d \geq 2$, there are nonclassical polynomials $P_{\varepsilon} : \mathbb{F}_p \rightarrow \mathbb{T}$, of degree at most d and phase order dividing $p^{\lceil d/(p-1) \rceil}$, such that*

$$\mathcal{M}_{\sigma}(f) \leq \prod_{\varepsilon \in \{0,1\}^d} \|fe(-P_{\varepsilon})\|_{U^d} \leq \max_{\varepsilon} \|fe(-P_{\varepsilon})\|_{U^d}^{2^d}. \quad (\text{IV.33.3})$$

Thus a lower bound $\mathcal{M}_{\sigma}(f) \geq \alpha$ is preserved as the same lower bound for the 2^d th power of a suitably twisted U^d norm.

Proof. Expand the square as in Proposition IV.33.1. Apply Theorem IV.30.2 to $h_d \sigma(h_1, \dots, h_{d-1})$, using $s = x$. This absorbs the character into a separate phase at each vertex, producing the mixed cube average of the functions $fe(-P_{\varepsilon})$. Gowers–Cauchy–Schwarz bounds its absolute value by the product in (IV.33.3); the second inequality is immediate. The left side is nonnegative, so this is the stated inequality. $\square$

The Gowers–Cauchy–Schwarz step is standard [1], and arbitrary vertex-phase absorption is already treated in the repository’s local refinement material [204]. The contribution here is the existence, explicit construction, and optimal universal phase alphabet of the phases in small characteristic, together with the exact obstruction analysis showing why the same statement cannot simply be transferred to arbitrary extension-field trace tensors.

Classical correlation is a different question. Failure of a classical total-derivative representation does not automatically mean that classical phases cannot provide the correlation required by an inverse theorem. These are different assertions. The distinction is central in the work of Berger–Sah–Sawhney–Tidor [186], and it is respected throughout this paper. In particular, the trace classification is a theorem about exact total derivatives, not a counterexample to an inverse theorem.

IV.34 Cube identities and phase invariance

Source 43, delivered Section 2.

IV.34.1 The finite-abelian normalization

Let G be a finite abelian group, written additively. For this section and Sections IV.35–IV.36, a multiadditive form takes values in $\mathbb{T}$. Define

$$\mathcal{E}_T(f) = \mathbb{E}_{\mathbf{h} \in G^{d-1}} |\mathbb{E}_x \partial_{\mathbf{h}} f(x) e(-T(\mathbf{h}, x))|^2. \quad (\text{IV.34.1})$$

In the Boolean specialization we use $T/2$ in this formula and retain the notation $\mathcal{E}_T$. This convention accounts for all factors of two.

Lemma IV.34.1 (Exact selected-cube identity). *For every complex-valued f and every multiadditive T ,*

$$\mathcal{E}_T(f) = \mathbb{E}_{x, h_1, \dots, h_d} \partial_{h_1} \cdots \partial_{h_d} f(x) e(-T(h_1, \dots, h_d)). \quad (\text{IV.34.2})$$

In particular the right side is real and nonnegative, even if T is not symmetric.

Proof. For fixed $\mathbf{h}$, put $g = \partial_{\mathbf{h}} f$. Expanding the square in (IV.34.1) and replacing the two base points by $x + y$ and x gives

$$|\mathbb{E}_x g(x) e(-T(\mathbf{h}, x))|^2 = \mathbb{E}_{x, y} g(x + y) \overline{g(x)} e(-T(\mathbf{h}, y)).$$

Average over $\mathbf{h}$. Multiplicative differences commute, so the product is the d -fold difference in (IV.34.2). $\square$

We define the Gowers quantity by

$$\|f\|_{U^d}^{2^d} = \mathbb{E}_{x, \mathbf{h} \in G^d} \partial_{\mathbf{h}} f(x). \quad (\text{IV.34.3})$$

Only this finite identity is needed below. Its nonnegativity follows from Lemma IV.34.1 with $T = 0$. For 1-bounded f , each selected Fourier coefficient has modulus at most one, so $0 \leq \mathcal{E}_T(f) \leq 1$.

Lemma IV.34.2 (Gauge invariance). *Suppose $P : G \rightarrow \mathbb{T}$ has constant d -fold derivative $S : G^d \rightarrow \mathbb{T}$. Then*

$$\mathcal{E}_{T+S}(fe(P)) = \mathcal{E}_T(f). \quad (\text{IV.34.4})$$

If $S = T$, then

$$\mathcal{E}_T(f) = \|fe(-P)\|_{U^d}^{2^d}. \quad (\text{IV.34.5})$$

Moreover the corresponding selected Fourier magnitudes agree for each individual $(d-1)$ -tuple of directions, before averaging.

Proof. Multiplicative differences distribute over products. Thus the extra factor in the cube identity for $fe(P)$ is $e(S)$, which cancels the extra tensor in the phase. For the pointwise assertion, the function $q_{\mathbf{h}}(x) = \Delta_{\mathbf{h}}P(x)$ satisfies $q_{\mathbf{h}}(x+y) - q_{\mathbf{h}}(x) = T(\mathbf{h}, y)$. Taking $x = 0$ shows $q_{\mathbf{h}}(x) = c_{\mathbf{h}} + T(\mathbf{h}, x)$ for a constant $c_{\mathbf{h}}$. Consequently

$$\mathbb{E}_x \partial_{\mathbf{h}}(fe(-P))(x) = e(-c_{\mathbf{h}}) \mathbb{E}_x \partial_{\mathbf{h}} f(x) e(-T(\mathbf{h}, x)).$$

Taking absolute values proves the assertion. □

Remark IV.34.3 (Amplitude and conjugation). No step in these identities divides by f or assumes $|f| = 1$. In particular they remain valid at zeros. This matters in the shear argument, where the new function is a product of two translates of f . The scale of $\mathcal{E}_T$ is 2^d : replacing f by cf multiplies the energy by $|c|^{2^d}$.

IV.35 Division-free finite-abelian orthogonality

Source 43, delivered Section 3.

IV.35.1 Alternating bicharacters and their dimensions

Let $B : G \times G \rightarrow \mathbb{T}$ be biadditive. Define

$$\kappa_B(a, b) = e(B(a, b) - B(b, a)), \quad R_B = \{a : \kappa_B(a, b) = 1 \text{ for every } b\}, \quad d_B = \sqrt{|G/R_B|}. \quad (\text{IV.35.1})$$

The pairing κ_B is alternating: $\kappa_B(a, a) = 1$, and $\kappa_B(a, b) = \kappa_B(b, a)^{-1}$. This formulation uses no division by two, including when G has even order.

Lemma IV.35.1 (Finite symplectic decomposition). *The number d_B is an integer. There is a subgroup $H \supseteq R_B$ such that $H = H^\perp$ for κ_B and $[G : H] = d_B$. More precisely,*

$$G/R_B \cong \bigoplus_j (C_{m_j} \oplus C_{m_j}), \quad d_B = \prod_j m_j, \quad (\text{IV.35.2})$$

with orthogonal nondegenerate alternating pairings on the summands.

Proof. Pass to $K = G/R_B$, where the pairing is nondegenerate. The map $K \rightarrow \widehat{K}$ it induces is injective and hence bijective. Choose $x \in K$ of maximum order m , which is also the exponent of K . The character $\kappa_B(x, \cdot)$ has order m by injectivity. Choose y such that $\kappa_B(x, y)$ is a primitive m th root of unity. Then y also has order m .

Pairing a relation $ax + by = 0$ with x and y shows that a and b are both zero modulo m . The subgroup $L = \langle x, y \rangle$ is therefore $C_m \oplus C_m$, and its restricted pairing is nondegenerate. For any $z \in K$, appropriate multiples of x, y subtract the two pairings with x, y ; hence $K = L \oplus L^\perp$. Induction on $|K|$ gives (IV.35.2). In each block, take the cyclic subgroup generated by its first basis vector. Their direct sum is self-annihilating. Its inverse image in G is the required H . $\square$

IV.35.2 The mixed orthogonality theorem

Use the normalized Hilbert space inner product $\langle g, h \rangle = \mathbb{E}_x g(x) \overline{h(x)}$, linear in its first argument. Define

$$\mathcal{Q}_B(g, h) = \mathbb{E}_a \left| \mathbb{E}_x g(x+a) \overline{h(x)} e(-B(a, x)) \right|^2. \quad (\text{IV.35.3})$$

Write $\mathcal{Q}_B(g) = \mathcal{Q}_B(g, g)$.

Theorem IV.35.2 (Sharp mixed orthogonality). *For every finite abelian G , every biadditive B , and all $g, h : G \rightarrow \mathbb{C}$,*

$$\mathcal{Q}_B(g, h) \leq d_B^{-1} \|g\|_2^2 \|h\|_2^2. \quad (\text{IV.35.4})$$

For each B there is a function g with $|g(x)| = 1$ at every x for which $\mathcal{Q}_B(g) = d_B^{-1}$. Thus the coefficient is sharp even in the one-function, pointwise-unit-modulus class.

Proof. Introduce unitaries

$$(U_a g)(x) = e(-B(a, x)) g(x+a). \quad (\text{IV.35.5})$$

Direct multiplication gives

$$U_a U_b = e(-B(b, a)) U_{a+b}, \quad U_a U_b = \kappa_B(a, b) U_b U_a. \quad (\text{IV.35.6})$$

Let M_g be the positive operator

$$M_g v = \mathbb{E}_a \langle v, U_a g \rangle U_a g.$$

It has trace $\|g\|_2^2$. The projective multiplication law shows that conjugation by each U_b merely permutes its summands, so M_g commutes with every U_b .

We claim that every nonzero subspace invariant under all U_b has dimension at least d_B . Choose H as in Lemma IV.35.1. The unitaries U_h , $h \in H$, commute, so they have a common eigenvector in any such subspace. If its joint eigenvalue function is $\lambda(h)$, then U_t transports it to the joint eigenvalue function $\kappa_B(h, t) \lambda(h)$. Two such functions agree precisely when the corresponding t 's differ by an element of $H^\perp = H$. Distinct joint eigenvalue functions belong to orthogonal subspaces. Representatives of G/H therefore give d_B independent vectors, proving the claim.

Each nonzero eigenspace of M_g is invariant under the U_b . Every positive eigenvalue consequently has multiplicity at least d_B , and

$$\|M_g\|_{\text{op}} \leq \frac{\text{tr } M_g}{d_B} = d_B^{-1} \|g\|_2^2.$$

Finally, $\mathcal{Q}_B(g, h) = \langle M_g h, h \rangle \leq \|M_g\|_{\text{op}} \|h\|_2^2$.

For sharpness, choose a common projective eigenvalue function λ for the commuting U_h , $h \in H$, and construct a λ -eigenvector of modulus one on every coset of H . Here is why this is possible. On a single coset, a joint eigenspace has dimension at most one, since translations by H determine all its coordinates from one coordinate. A nonzero eigenvector has constant, nonzero modulus on

that coset. There are $|H|$ distinct joint eigenvalue functions on each coset. The ratio of any two eigenvalue functions with the same projective multiplication law is an ordinary character of H . There are exactly $|H|$ such ratios. Thus every coset has the same complete set of joint eigenvalue functions. Normalize and assemble their λ -vectors to obtain a global g with $|g| = 1$.

For $a \in H$, $U_a g = \lambda(a)g$, so $|\langle U_a g, g \rangle| = 1$. For $a \notin H$, choose $h \in H$ with $\kappa_B(h, a) \neq 1$. The vectors g and $U_a g$ have different U_h -eigenvalues and are orthogonal. It follows that $\mathcal{Q}_B(g) = |H|/|G| = d_B^{-1}$. $\square$

Corollary IV.35.3 (Self-cube form). *For every g ,*

$$\mathbb{E}_{a,b,x} \partial_a \partial_b g(x) e(-B(a, b)) = \mathcal{Q}_B(g) \leq d_B^{-1} \|g\|_2^4. \quad (\text{IV.35.7})$$

Proof. Expand the square in (IV.35.3), replace the two base points by $x + b$ and x , and use biadditivity. Then apply Theorem IV.35.2. $\square$

IV.35.3 Fields, elementary divisors, and a composite example

If $G = \mathbb{F}_p^n$ and B is represented by a matrix M through $B(a, b) = a^\top M b / p$, then

$$d_B = p^{\text{rank}(M - M^\top)/2}. \quad (\text{IV.35.8})$$

An alternating bilinear form has even rank also in characteristic two; this is the vector-space case of Lemma IV.35.1.

For $G = (\mathbb{Z}/m\mathbb{Z})^n$, let M be an integer matrix and $B(a, b) = a^\top M b / m \pmod{1}$. If $s_1, \dots, s_n$ are the Smith invariants of the integer matrix $M - M^\top$, including zeros, then

$$d_B = \left(\prod_{i=1}^n \frac{m}{\gcd(m, s_i)} \right)^{1/2}, \quad \gcd(m, 0) = m. \quad (\text{IV.35.9})$$

Indeed unimodular row and column operations remain invertible modulo m , and multiplication by s_i has kernel of size $\gcd(m, s_i)$. The square root is an integer by the preceding theorem.

For example, $B(a, b) = s a_1 b_2 / m$ on $(\mathbb{Z}/m\mathbb{Z})^2$ has $d_B = m / \gcd(m, s)$. The constant function already gives $\mathcal{Q}_B(1) = \gcd(m, s) / m$. At $m = 12$, the factors for $s = 1, 2, 3, 4, 6$ are respectively $1/12, 1/6, 1/4, 1/3, 1/2$. A rank computation over only one residue field cannot distinguish this arithmetic profile.

IV.36 Sharp symmetry thresholds on finite abelian groups

Source 43, delivered Section 4.

IV.36.1 Alternating slices

Let $T : G^{k+1} \rightarrow \mathbb{T}$ be multiadditive, with $k \geq 1$. Fix a slot $i \in \{1, \dots, k\}$, and freeze the other $k - 1$ derivative directions at $\mathbf{z}$. Let $B_{i,\mathbf{z}}(a, b)$ be the bilinear form obtained by putting a in slot i and b in the last slot of T .

Theorem IV.36.1 (Amplitude-sensitive slice profile). *For every such slot,*

$$\mathcal{E}_T(f) = \mathbb{E}_z \mathcal{Q}_{B_{i,z}}(\partial_z f) \leq \mathbb{E}_z d_{B_{i,z}}^{-1} \|\partial_z f\|_2^4. \quad (\text{IV.36.1})$$

For 1-bounded f the final norm factors may be omitted.

Proof. Use Lemma IV.34.1, freeze the other directions, and use commutation of multiplicative differences. The inner average is exactly (IV.35.7). Each frozen derivative is 1-bounded when f is. $\square$

Lemma IV.36.2 (Support of a multiadditive map). *Let P be a nontrivial finite abelian p -group and let $F : P^r \rightarrow A$ be a nonzero multiadditive map to an abelian group. Then*

$$\mathbb{P}(F(z_1, \dots, z_r) \neq 0) \geq (1 - 1/p)^r. \quad (\text{IV.36.2})$$

For $r = 0$, a nonzero constant has support probability one.

Proof. A nonzero homomorphism out of P has image of size at least p ; its kernel therefore occupies at most $1/p$ of P . For the induction, view F as a map of the first $r - 1$ variables into $\text{Hom}(P, A)$. This multiadditive map is nonzero. On the at least $(1 - 1/p)^{r-1}$ fraction where its value is a nonzero homomorphism, the last argument is nonzero in the image with probability at least $1 - 1/p$. $\square$

Corollary IV.36.3 (Prime-power symmetry gap). *If G is a finite abelian p -group and T is not invariant under interchanging slot i and the last slot, then*

$$\mathcal{E}_T(f) \leq 1 - (1 - 1/p)^k. \quad (\text{IV.36.3})$$

Proof. The alternating bicharacter of $B_{i,z}$ is a nonzero $(k - 1)$ -variable multiadditive map into the group of bicharacters. It is nontrivial on a fraction at least $(1 - 1/p)^{k-1}$ of the tuples. On those tuples its nondegenerate quotient has order at least p^2 by Lemma IV.35.1; hence $d_{B_{i,z}}^{-1} \leq 1/p$. On the other tuples the factor is one. The total loss is at least $(1 - 1/p)^k$. $\square$

IV.36.2 Primary factorization and the exact group threshold

Write $G = \bigoplus_p G_p$. Multiadditivity implies

$$T(h_1, \dots, h_{k+1}) = \sum_p T_p(h_{1,p}, \dots, h_{k+1,p}). \quad (\text{IV.36.4})$$

To see this, expand each argument into primary components. A term containing inputs of two different prime-power orders is annihilated by two coprime integers and is therefore zero in $\mathbb{T}$.

Let $\mathcal{P}_i$ be the primes for which the slot- i alternating defect of T_p is nonzero. The radical indices multiply, and averaging over the independent primary coordinates yields

$$\mathbb{E}_z d_{B_{i,z}}^{-1} = \prod_p \mathbb{E}_{z_p} d_{B_{i,z_p}}^{-1} \leq \prod_{p \in \mathcal{P}_i} [1 - (1 - 1/p)^k]. \quad (\text{IV.36.5})$$

The energy itself need not factor for an arbitrary f . It is the *upper rank profile* in (IV.36.5) that factors.

Theorem IV.36.4 (Exact universal symmetry threshold). *Suppose G is not cyclic, and let $p_*(G)$ be the least prime p for which G has a quotient isomorphic to $\mathbb{F}_p^2$. Then*

$$\mathcal{E}_T(f) > 1 - (1 - 1/p_*(G))^k \implies T \text{ is fully symmetric.} \quad (\text{IV.36.6})$$

The threshold is attained by a nonsymmetric T and $f = 1$, even when T is symmetric in its first k slots. If G is cyclic, every multiadditive $T : G^{k+1} \rightarrow \mathbb{T}$ is already symmetric.

Proof. On a cyclic group with generator g , multilinearity gives

$$T(a_1g, \dots, a_{k+1}g) = a_1 \cdots a_{k+1}T(g, \dots, g).$$

This expression is symmetric. Thus a nonzero alternating defect must occur on a noncyclic primary component. Such a component has a quotient $\mathbb{F}_p^2$, so $p \geq p_*(G)$. The function $1 - (1 - 1/p)^k$ decreases as p increases. Corollary IV.36.3 and (IV.36.5) now bound the energy of any nonzero slot defect by the right threshold. The transpositions of the last slot with each other slot generate all permutations. Absence of these defects is full symmetry.

For sharpness, choose independent linear forms u, v on a quotient $\mathbb{F}_p^2$ with $p = p_*(G)$, pull them back to G , and set

$$T(h_1, \dots, h_k, y) = \frac{u(h_1) \cdots u(h_k)v(y)}{p} \pmod{1}. \quad (\text{IV.36.7})$$

It is nonsymmetric: choose an input where all the $u(h_j)$ and $v(y)$ are one, while $u(y) = 0$ and one of the exchanged $v(h_j)$ is zero. For $f = 1$ a selected coefficient is one precisely when $u(h_1) \cdots u(h_k) = 0$, and is zero otherwise. This has probability $1 - (1 - 1/p)^k$. The first k slots are symmetric by construction. $\square$

[write] **7 October 2026; sources 40–46.** For $G = \mathbb{F}_p^n$, p odd, this is source 36's Theorem F.41.1; here it holds on every finite abelian group, characteristic two included, with $p_*(G)$ in place of p , and the orthogonality factor d_B^{-1} of Theorem IV.35.2 extends source 36's Theorem IV.17.1 (credited by the source). This advances source 36's composite-modulus question (item IV.9).

Remark IV.36.5 (Two independent kinds of sharpness). Taking sums of (IV.36.7) on distinct primary components and $f = 1$ attains the product in (IV.36.5). For a fixed bilinear form, the exact factor d_B^{-1} is attained by the construction in Theorem IV.35.2. Neither claim requires that an arbitrary multilinear rank profile be simultaneously attainable by one function; such a claim would be stronger and is not made here.

IV.37 Boolean nonclassical integration, explicitly

Source 43, delivered Section 5.

This section supplies a self-contained Boolean form of the existing nonclassical integration theorem. Its role is to specify exactly which algebraic conditions the new spectral estimates must certify.

Definition IV.37.1. Write $\text{SymMult}^d(V)$ for the space of symmetric d -linear maps $V^d \rightarrow \mathbb{F}_2$. For $d \geq 3$, define $\text{nCSM}^d(V)$ to be the subspace satisfying

$$T(a, a, b, \mathbf{z}) = T(a, b, b, \mathbf{z}) \quad \text{for every } a, b \in V, \mathbf{z} \in V^{d-3}. \quad (\text{IV.37.1})$$

For $d = 2$ put $\text{nCSM}^2(V) = \text{SymMult}^2(V)$. The notation denotes a space of symmetric *multilinear maps*; no identification with symmetric tensor coinvariants in characteristic two is being assumed.

Theorem IV.37.2 (Constructive Boolean integration). *A d -linear form $T : V^d \rightarrow \mathbb{F}_2$, $d \geq 2$, has a nonclassical primitive of degree at most d if and only if $T \in \text{nCSM}^d(V)$. Fix coordinates $x_1, \dots, x_n$ with basis $e_1, \dots, e_n$. For each nonempty $S \subseteq \{1, \dots, n\}$ with $|S| \leq d$, let c_S be the value of T on any length- d basis tuple whose set of indices is exactly S . This value is well defined, and a primitive is*

$$P_T(x) = \sum_{\substack{\emptyset \neq S \subseteq [n] \\ |S| \leq d}} \frac{c_S \prod_{i \in S} |x_i|}{2^{d-|S|+1}} \pmod{1}, \quad |x_i| \in \{0, 1\}. \quad (\text{IV.37.2})$$

In particular,

$$\dim_{\mathbb{F}_2} \text{nCSM}^d(V) = \sum_{s=1}^{\min(n,d)} \binom{n}{s}. \quad (\text{IV.37.3})$$

Proof. Let S_a be translation by a on additive functions. Because $2a = 0$, $S_a^2 = I$, and consequently

$$\Delta_a^2 = (S_a - I)^2 = -2\Delta_a, \quad \Delta_a^2 \Delta_b = \Delta_a \Delta_b^2. \quad (\text{IV.37.4})$$

The second identity follows since both sides equal $-2\Delta_a \Delta_b$. Additive differences commute. A constant top derivative is therefore symmetric and satisfies (IV.37.1). This proves necessity.

Now assume $T \in \text{nCSM}^d(V)$. By symmetry, a coefficient on basis vectors depends only on their multiplicities. If two distinct indices i, j are present and i occurs at least twice, the identity (IV.37.1) changes one occurrence of i into one of j without changing the coefficient. For any fixed support S , all positive multiplicity vectors of total d are connected by these moves: move excess occurrences into a chosen index of S , then redistribute to the desired multiplicities. Thus c_S is well defined. When $d = 2$ this assertion requires no moves.

It remains to check the proposed monomials, including their degree. For a set S of size s , put

$$P_S(x) = 2^{-(d-s+1)} \prod_{i \in S} |x_i| \pmod{1}.$$

On a single Boolean coordinate, repeated application of (IV.37.4) gives $\Delta_{e_i}^t = (-2)^{t-1} \Delta_{e_i}$ for $t \geq 1$. Also

$$\Delta_h = \prod_{i: h_i=1} (I + \Delta_{e_i}) - I. \quad (\text{IV.37.5})$$

Expanding any $(d+1)$ differences by (IV.37.5) produces products of at least $d+1$ coordinate differences. If any used coordinate lies outside S , the term is zero. Otherwise, if $r \leq s$ distinct coordinates occur among $t \geq d+1$ differences, the numerator is divisible by 2^{t-r} . Since $t-r \geq d-s+1$, its quotient by the denominator is integer-valued. It vanishes in $\mathbb{T}$. Thus P_S has degree at most d .

Apply exactly d coordinate differences to P_S . If their support contains an index outside S , the answer is zero. If their support is a proper subset $R \subsetneq S$, of size $r < s$, the numerator is divisible by 2^{d-r} , giving an integer quotient because $d-r \geq d-s+1$. If their support is exactly S , the result is $\pm 1/2 \pmod{1} = 1/2$, independent of the base point. Hence the top derivative on a basis tuple has value $1/2$ exactly when that tuple's support is S .

For completeness, the top derivative of any degree-at-most- d polynomial is multiadditive. The identity $\Delta_{a+b} = \Delta_a + \Delta_b + \Delta_a \Delta_b$ shows that its failure of additivity is a $(d+1)$ -fold derivative and therefore zero. It is independent of the base point for the same reason. Agreement on basis tuples now proves that the top derivative of (IV.37.2) is $T/2$ everywhere. The support coefficients are independent, as the same basis-tuple calculation isolates each one. Counting them proves (IV.37.3). $\square$

Example IV.37.3 (Why fourth and eighth roots are unavoidable). On $\mathbb{F}_2$, $P(x) = |x|/4$ has second derivative $ab/2$, and $P(x) = |x|/8$ has third derivative $abc/2$. Thus symmetric diagonal terms integrate perfectly, but not generally into $\mathbb{F}_2$ -valued classical phases. Already the quadratic diagonal requires fourth roots of unity. This is not an obstruction to the value group $\mathbb{T}$ used throughout the manuscript.

Remark IV.37.4 (Dependence on coordinates). The primitive in (IV.37.2) is a chosen section of the top-derivative map, not a coordinate-free canonical polynomial. Two primitives differ by a polynomial of degree at most $d - 1$. Their use in phase removal yields the same selected energy because such a difference has zero top derivative.

IV.38 The shear identity and sharp Boolean certification

Source 43, delivered Section 6.

IV.38.1 A bilinear form hidden in a repeated direction

We return to a symmetric d -linear form $T : V^d \rightarrow \mathbb{F}_2$, $d \geq 3$. For $\mathbf{z} \in V^{d-3}$ define $\Omega_{\mathbf{z}}$ by (F.48.8).

Lemma IV.38.1. *Each $\Omega_{\mathbf{z}}$ is an alternating bilinear form. The map $\mathbf{z} \mapsto \Omega_{\mathbf{z}}$ is multilinear with values in $\text{Alt}^2(V)$.*

Proof. In characteristic two, symmetry cancels the two cross terms when a is replaced by $a + a'$ in $T(a, a, b, \mathbf{z})$; hence this expression is linear in a . The same argument applies to the repeated b . The remaining occurrences are already linear. Finally $\Omega_{\mathbf{z}}(a, a) = 2T(a, a, a, \mathbf{z}) = 0$, and multilinearity in the frozen variables is immediate. $\square$

For a function f , write $\tau_w f(x) = f(x + w)$. The elementary identity responsible for the stronger estimates is

$$\partial_a \partial_{a+w} f = \overline{\partial_a (f \tau_w f)}. \quad (\text{IV.38.1})$$

Indeed both sides at x equal $f(x)f(x+w)\overline{f(x+a)f(x+a+w)}$. This identity has no exceptional set where f vanishes.

Theorem IV.38.2 (Exact shear and repeated-variable rank profile). *Put*

$$g_{\mathbf{z},w} = \partial_{\mathbf{z}}(f \tau_w f), \quad (\text{IV.38.2})$$

$$B_{\mathbf{z},w}(a, b) = T(a, a, b, \mathbf{z}) + T(a, b, \mathbf{z}, w). \quad (\text{IV.38.3})$$

The form $B_{\mathbf{z},w}$ is bilinear and $B_{\mathbf{z},w}(a, b) + B_{\mathbf{z},w}(b, a) = \Omega_{\mathbf{z}}(a, b)$. In the notation where Boolean bilinear forms are embedded in $\mathbb{T}$ by division by two, one has the exact identity

$$\mathcal{E}_T(f) = \mathbb{E}_{\mathbf{z},w} \mathcal{Q}_{B_{\mathbf{z},w}}(g_{\mathbf{z},w}). \quad (\text{IV.38.4})$$

Consequently, for arbitrary complex f ,

$$\mathcal{E}_T(f) \leq \mathbb{E}_{\mathbf{z},w} 2^{-\text{rank } \Omega_{\mathbf{z}}/2} \|g_{\mathbf{z},w}\|_2^4. \quad (\text{IV.38.5})$$

For 1-bounded f this gives

$$\mathcal{E}_T(f) \leq \mathbb{E}_{\mathbf{z}} 2^{-\text{rank } \Omega_{\mathbf{z}}/2}. \quad (\text{IV.38.6})$$

Proof. Write the cube variables in Lemma IV.34.1 as a, b, z, c , and make the invertible substitution $c = a + w$. Multilinearity gives

$$T(a, b, z, a + w) = T(a, a, b, z) + T(a, b, z, w) = B_{z,w}(a, b).$$

The first term is bilinear in a, b by the cancellation used in Lemma IV.38.1; the second term is bilinear and symmetric in a, b . Its alternating part cancels.

Apply the remaining differences $\partial_b \partial_z$ to (IV.38.1). They commute with complex conjugation. The resulting inner average over a, b, x is the complex conjugate of $\mathcal{Q}_{B_{z,w}}(g_{z,w})$, by Corollary IV.35.3. That quantity is real and nonnegative, so its conjugate is itself. This proves (IV.38.4). The orthogonality factor is $2^{-\text{rank } \Omega_z/2}$ by (IV.35.8). Applying Theorem IV.35.2 gives (IV.38.5). If f is 1-bounded, so is $g_{z,w}$, and the norm factors are at most one. $\square$

Corollary IV.38.3 (A strict gap for the additional obstruction). *If T is symmetric but not in $\text{nCSM}^d(V)$, then*

$$\mathcal{E}_T(f) \leq 1 - 2^{2-d}. \quad (\text{IV.38.7})$$

Proof. The multilinear map $z \mapsto \Omega_z$ is nonzero. Lemma IV.36.2 says that its support has measure at least $2^{-(d-3)}$. A nonzero alternating form has rank at least two, so its factor in (IV.38.6) is at most $1/2$. The average loses at least $2^{-(d-3)}/2 = 2^{2-d}$ from one. $\square$

IV.38.2 The sharp all-degree theorem

Theorem IV.38.4 (Sharp Boolean phase-integration criterion). *Let $d \geq 2$ and let $T : V^d \rightarrow \mathbb{F}_2$ be multilinear, without an assumed symmetry. If*

$$\mathcal{E}_T(f) > 1 - 2^{1-d}, \quad (\text{IV.38.8})$$

then $T \in \text{nCSM}^d(V)$. The polynomial P_T in (IV.37.2) exists and satisfies

$$\mathcal{E}_T(f) = \|fe(-P_T)\|_{U^d}^{2d}. \quad (\text{IV.38.9})$$

The threshold cannot be lowered in any dimension $n \geq 2$, even when T is assumed symmetric in its first $d - 1$ variables.

Proof. Theorem IV.36.4, with $p = 2$ and $k = d - 1$, forces full symmetry. For $d = 2$ this already gives integrability. For $d \geq 3$, failure of integrability would imply (IV.38.7). But $1 - 2^{2-d} < 1 - 2^{1-d}$, contradicting (IV.38.8). Integration follows from Theorem IV.37.2, and the exact identity follows from Lemma IV.34.2. The nonsymmetric endpoint in (IV.36.7), using two independent Boolean linear forms, has precisely energy $1 - 2^{1-d}$ and cannot be a top derivative. This proves sharpness. $\square$

Remark IV.38.5 (The two thresholds should not be conflated). The universal threshold for a prescribed multilinear selector is $1 - 2^{1-d}$. For a tensor already known to be symmetric, the sufficient threshold $1 - 2^{2-d}$ is smaller. The latter is proved optimal here only for $d = 3$; Section IV.41 shows it is not optimal for quartics in dimension two. In dimension one all symmetric multilinear tensors integrate, so the dimension-independent sharpness examples necessarily use at least two coordinates.

IV.38.3 Corrupted selectors and finite margins

Proposition IV.38.6 (Robust certification). *Let $\rho : V^{d-1} \rightarrow V^*$ be an arbitrary selector, and suppose it agrees outside a fraction η of tuples with $\sigma(\mathbf{h}) = T(\mathbf{h}, \cdot)$ for a multilinear T . Define $\mathcal{E}_\rho(f)$ by the corresponding selected Fourier energy. Then*

$$|\mathcal{E}_\rho(f) - \mathcal{E}_T(f)| \leq \eta. \quad (\text{IV.38.10})$$

If a certified lower bound L on $\mathcal{E}_\rho(f)$ satisfies $L - \eta > 1 - 2^{1-d}$, then T integrates and

$$\|fe(-P_T)\|_{U^d}^{2^d} \geq L - \eta. \quad (\text{IV.38.11})$$

If T is known symmetric, $1 - 2^{2-d}$ may replace the threshold for $d \geq 3$. In the symmetric cubic case, every $L - \eta > 0$ also gives an integrable restriction of codimension at most $\lfloor \log_2(1/(L - \eta)) \rfloor$.

Proof. Each squared selected Fourier coefficient lies in $[0, 1]$. The summands agree off the exceptional set, so their average difference is at most its measure. Apply the preceding theorems; the cubic assertion uses Theorem IV.39.4 below. $\square$

The hypothesis supplies a genuine multilinear model and a certified energy bound. It does not assert that noisy samples automatically recover such a model. Numerical error in a computed lower bound must be subtracted before using the strict margin.

IV.39 Exact cubic obstruction ranks and optimal repair

Source 43, delivered Section 7.

IV.39.1 The amplitude-sensitive inequality

Theorem IV.39.1 (Cubic obstruction bound). *Let $T : V^3 \rightarrow \mathbb{F}_2$ be symmetric, and write $\text{rank } \Omega_T = 2m$. For every $f : V \rightarrow \mathbb{C}$,*

$$\mathcal{E}_T(f) \leq 2^{-m} \left\| |f|^2 \right\|_{U^2}^4. \quad (\text{IV.39.1})$$

In particular, if $|f| \leq 1$ and $\mu = \mathbb{E}_x |f(x)|^2$, then

$$\mathcal{E}_T(f) \leq 2^{-m} \mu^3 \leq 2^{-m}. \quad (\text{IV.39.2})$$

Proof. In Theorem IV.38.2, there are no frozen variables $\mathbf{z}$. The rank factor is constant, and $g_w = f\tau_w f$. Thus the remaining factor is

$$\mathbb{E}_w \|f\tau_w f\|_2^4 = \mathbb{E}_w \left(\mathbb{E}_x |f(x)|^2 |f(x+w)|^2 \right)^2 = \left\| |f|^2 \right\|_{U^2}^4.$$

For $a = |f|^2 \in [0, 1]$, put $c(w) = \mathbb{E}_x a(x)a(x+w)$. Then $0 \leq c(w) \leq \mu$ and $\mathbb{E}_w c(w) = \mu^2$. Consequently $\mathbb{E}_w c(w)^2 \leq \mu \mathbb{E}_w c(w) = \mu^3$. $\square$

IV.39.2 A normal form modulo integrable tensors

For two linear forms u, v , define the symmetric trilinear tensor

$$C_{u,v}(a, b, c) = u(a)u(b)v(c) + u(a)v(b)u(c) + v(a)u(b)u(c). \quad (\text{IV.39.3})$$

A direct substitution gives

$$\Omega_{C_{u,v}}(a, b) = u(a)v(b) + v(a)u(b). \quad (\text{IV.39.4})$$

Proposition IV.39.2 (The cubic obstruction quotient). *The linear map*

$$\Omega : \text{SymMult}^3(V) \longrightarrow \text{Alt}^2(V), \quad T \longmapsto \Omega_T,$$

is surjective and has kernel $\text{nCSM}^3(V)$. It induces a $\text{GL}(V)$ -equivariant isomorphism

$$\text{SymMult}^3(V) / \text{nCSM}^3(V) \cong \text{Alt}^2(V). \quad (\text{IV.39.5})$$

If $\text{rank } \Omega_T = 2m$, there are independent forms $u_1, v_1, \dots, u_m, v_m$ such that

$$T = S + C_m, \quad S \in \text{nCSM}^3(V), \quad C_m = \sum_{j=1}^m C_{u_j, v_j}. \quad (\text{IV.39.6})$$

Proof. The kernel assertion is precisely the repeated-variable identity. Every alternating bilinear form has a symplectic decomposition $\sum_j (u_j \otimes v_j + v_j \otimes u_j)$; this is the vector-space proof in Lemma IV.35.1, with additional radical coordinates. Formula (IV.39.4) gives a preimage of each block and proves (IV.39.6). The definition of Ω_T commutes with every linear change of coordinates, proving equivariance. $\square$

Lemma IV.39.3 (Canonical cubic energy). *For the tensor C_m in (IV.39.6),*

$$\mathcal{E}_{C_m}(1) = 2^{-m}. \quad (\text{IV.39.7})$$

Proof. For a constant function, the energy is the probability that the selected linear functional $C_m(a, b, \cdot)$ vanishes. Different symplectic coordinate pairs are independent, so it suffices to count one pair. Write $a = (a_1, a_2)$ and $b = (b_1, b_2)$. The functional is

$$y \longmapsto a_1 b_1 y_2 + (a_1 b_2 + a_2 b_1) y_1.$$

For $a = (0, 0)$, all four b 's work. For $a = (0, 1)$, two work. For each of $a = (1, 0)$ and $(1, 1)$, exactly one works. There are $4 + 2 + 1 + 1 = 8$ successful pairs out of sixteen, giving $1/2$ per block. Radical coordinates impose no condition. Multiplication over m blocks proves the result. $\square$

IV.39.3 The exact extremal and repair theorem

Theorem IV.39.4 (Exact cubic energy and optimal codimension). *For every symmetric trilinear $T : V^3 \rightarrow \mathbb{F}_2$ with $\text{rank } \Omega_T = 2m$,*

$$\mathcal{M}(T) = 2^{-m}, \quad \min_{\substack{U \leq V \\ T|_{U^3} \in \text{nCSM}^3(U)}} \text{codim } U = m. \quad (\text{IV.39.8})$$

The maximum is attained by a nonclassical cubic phase of modulus one. An optimal subspace and an explicit primitive on it are constructible by finite linear algebra. Hence, for $0 < a \leq \mathcal{E}_T(f)$,

$$m \leq \lfloor \log_2(1/a) \rfloor. \quad (\text{IV.39.9})$$

The coefficient one in this logarithmic repair law is sharp.

Proof. The upper bound is Theorem IV.39.1. Write $T = S + C_m$ by Proposition IV.39.2, and integrate S to a cubic P_S . Gauge invariance and Lemma IV.39.3 give

$$\mathcal{E}_T(e(P_S)) = \mathcal{E}_{C_m}(1) = 2^{-m}.$$

This proves the exact extremum and its phase attainment.

By Theorem IV.37.2, $T|_{U^3}$ integrates precisely when $\Omega_T|_{U \times U} = 0$, that is, when U is isotropic. If $R = \text{rad } \Omega_T$, then $\dim R = n - 2m$. The image of an isotropic U in the nondegenerate symplectic quotient V/R has dimension at most m . Indeed in a nondegenerate alternating space of dimension $2m$, an isotropic subspace L satisfies $L \subseteq L^\perp$ and $\dim L + \dim L^\perp = 2m$. Thus

$$\dim U \leq \dim R + m = n - m.$$

Equality is obtained by adjoining the radical to a Lagrangian of the quotient. One can take $U = \bigcap_{j=1}^m \ker u_j$ in the normal form. Integrate the restriction using (IV.37.2) on a basis of U . Finally $a \leq 2^{-m}$ gives (IV.39.9). For every m , the example $T = C_m$, $f = 1$ has $a = 2^{-m}$ and requires exactly codimension m , so no smaller universal coefficient can replace one. $\square$

[write] **7 October 2026; sources 40–46.** This is source 46's Theorem IV.27.1 for one function, proved independently the same day; source 46's statement for eight functions contains the value $2^{-\text{rank } \Omega_T/2}$, and both identify the optimal codimension $\text{rank } \Omega_T/2$.

Corollary IV.39.5 (A low-partition-rank global correction). *There is $S \in \text{nCSM}^3(V)$ with*

$$\text{prank}(T - S) \leq 2m \leq 2\lfloor \log_2(1/a) \rfloor \quad \text{whenever } \mathcal{E}_T(f) \geq a > 0. \quad (\text{IV.39.10})$$

Here partition rank means the least number of summands, each factoring into a function of one nonempty set of variables and a function of its complementary set.

Proof. Use the normal form, and write each block as

$$C_{u,v}(a, b, c) = u(a)(u(b)v(c) + v(b)u(c)) + v(a)u(b)u(c).$$

This gives two partition-rank-one summands per block. The assertion is an upper bound; optimality of this partition-rank coefficient is not asserted. $\square$

IV.39.4 Concentration, tensor products, and what stability is available

The amplitude inequality also has concentrated extremizers. Choose an optimal isotropic subspace U of codimension m and a primitive P_U of $T|_{U^3}$. Define $f = 1_U e(P_U)$ on U and zero elsewhere. A nonzero cube contribution requires all its vertices in U ; this is equivalent to $x, a, b, c \in U$. Every such contribution is one after the tensor twist. With $\mu = |U|/|V| = 2^{-m}$,

$$\mathcal{E}_T(f) = \mu^4 = 2^{-m} \mu^3. \quad (\text{IV.39.11})$$

Thus the mass factor in (IV.39.2) can also be attained.

For direct sums of symmetric cubic tensors, alternating ranks add. The exact theorem therefore gives the multiplicative law

$$\mathcal{M}(T_1 \oplus T_2) = \mathcal{M}(T_1) \mathcal{M}(T_2). \quad (\text{IV.39.12})$$

This identity is a consequence of the proved extremal formula, not an assumption that arbitrary optimizing functions factor.

A first quantitative stability consequence concerns amplitudes. If $\mathcal{E}_T(f) \geq (1 - \varepsilon)2^{-m}$ with $0 \leq \varepsilon \leq 1$, then

$$\mathbb{E}_x(1 - |f(x)|^2) \leq 1 - (1 - \varepsilon)^{1/3}. \quad (\text{IV.39.13})$$

This follows immediately from (IV.39.2). It does not classify the phases of near-extremizers. A phase-stability claim would require information about simultaneous equality in the projective spectral inequalities and is left for further research.

IV.40 Mixed correlations: an optimal fourth power and a smaller repair bound

Source 43, delivered Section 8.

The one-function cube structure is stronger than an arbitrary mixed correlation. Nevertheless the *mixed* orthogonality theorem also improves a published local estimate without changing its hypotheses.

IV.40.1 The triangle correlation

For a finite abelian group and a biadditive $B : G^2 \rightarrow \mathbb{T}$, put

$$L_B(u, v, w) = \mathbb{E}_{a,b} u(a)v(b)w(a+b)e(B(a, b)). \quad (\text{IV.40.1})$$

Theorem IV.40.1 (Sharp fourth-power bilinear estimate). *For arbitrary complex u, v, w ,*

$$|L_B(u, v, w)| \leq d_B^{-1/2} \|u\|_2 \|v\|_2 \|w\|_2. \quad (\text{IV.40.2})$$

Consequently, for 1-bounded functions,

$$|L_B(u, v, w)|^4 \leq \mathbb{E}_{a,b} e(B(a, b) - B(b, a)) = d_B^{-2}. \quad (\text{IV.40.3})$$

The exponent four in (IV.40.3) is optimal as a uniform assertion, already on $\mathbb{F}_2^2$.

Proof. Write $B^t(a, b) = B(b, a)$, and use the operators associated to B^t . Replacing a by $x - b$ gives

$$\mathbb{E}_a u(a)w(a+b)e(B(a, b)) = e(-B(b, b)) \left\langle U_{-b}^{B^t} u, \bar{w} \right\rangle.$$

Cauchy–Schwarz in b , followed by Theorem IV.35.2, yields

$$|L_B|^2 \leq \|v\|_2^2 \mathcal{Q}_{B^t}(u, \bar{w}) \leq d_B^{-1} \|u\|_2^2 \|v\|_2^2 \|w\|_2^2.$$

The radical indices of B and B^t agree. Averaging the alternating character in its second variable is one when the first variable is in R_B and zero otherwise. Hence its bias is $|R_B|/|G| = d_B^{-2}$, proving (IV.40.3).

For sharpness, take $G = \mathbb{F}_2^2$ and $B(a, b) = a_1 b_2 / 2$. Then $d_B = 2$. Let $\zeta = e^{\pi i / 4}$ and choose

$$u(a) = 1, \quad w(x) = i^{x_1}, \quad v(b) = \begin{cases} \zeta^{-1}, & b_2 = 0, \\ \zeta(-1)^{b_1}, & b_2 = 1. \end{cases} \quad (\text{IV.40.4})$$

The inner average over a is $\zeta/\sqrt{2}$ when $b_2 = 0$, and $(-1)^{b_1}\zeta^{-1}/\sqrt{2}$ when $b_2 = 1$. Multiplying by $v(b)$ makes it $1/\sqrt{2}$ in both cases. Thus $|L_B| = 2^{-1/2}$ and $|L_B|^4 = 1/4 = d_B^{-2}$. Any exponent smaller than four would fail for this example. $\square$

Tidor's Lemma 4.1 [150] gives the eighth-power version of this estimate over finite fields. The fourth-power statement above is therefore a strict local refinement of that displayed estimate, as well as a finite-abelian extension. It makes no claim about priority among all possible earlier formulations of projective orthogonality.

IV.40.2 A transposition defect under seven-function correlation

For $T : V^3 \rightarrow \mathbb{F}_p$ and 1-bounded $b_1, \dots, b_7$, write

$$\Gamma_T(b_1, \dots, b_7) = \left| \mathbb{E}_{x,y,z} b_1(x) b_2(y) b_3(z) b_4(x+y) \cdot b_5(x+z) b_6(y+z) b_7(x+y+z) e_p(T(x,y,z)) \right|, \quad (\text{IV.40.5})$$

where $e_p(t) = e(t/p)$. Let $T_{(12)}(x,y,z) = T(y,x,z)$.

Proposition IV.40.2 (Fourth-power control of each transposition). *If $\Gamma_T \geq \delta > 0$, then*

$$\mathbb{E}_{x,y,z} e_p(T(x,y,z) - T(y,x,z)) \geq \delta^4. \quad (\text{IV.40.6})$$

The same conclusion holds for each transposition of the three slots. Equivalently, the analytic rank of each such defect is at most $4 \log_p(1/\delta)$, where analytic rank is minus the base- p logarithm of the bias.

Proof. Freeze z . The inner expression is a triangle correlation in x, y : combine $b_1(x)b_5(x+z)$, $b_2(y)b_6(y+z)$, and $b_4(x+y)b_7(x+y+z)$ into the three functions. If K_z is the alternating part of the bilinear form $T(\cdot, \cdot, z)$, Theorem IV.40.1 gives the upper bound $p^{-\text{rank } K_z/4}$. The triangle inequality and concavity of the fourth root give

$$\delta \leq \mathbb{E}_z p^{-\text{rank } K_z/4} \leq \left(\mathbb{E}_z p^{-\text{rank } K_z} \right)^{1/4}.$$

The average on the right is precisely the bias in (IV.40.6). Permuting the variables permutes the seven allowed arguments and proves the other cases. $\square$

This proposition only controls analytic rank. It does not silently replace analytic rank by partition rank for general trilinear tensors. Such a conversion is a separate input in a full inverse argument.

IV.40.3 Refining the symmetric trilinear repair lemma

Theorem IV.40.3 (Seven-function repair with coefficient two). *Let $T : V^3 \rightarrow \mathbb{F}_2$ be symmetric, and let $\text{rank } \Omega_T = 2m$. For every seven 1-bounded functions,*

$$\Gamma_T(b_1, \dots, b_7) \leq 2^{-m/2}. \quad (\text{IV.40.7})$$

Thus $\Gamma_T \geq \delta > 0$ implies that T integrates on a subspace of codimension at most

$$\lfloor 2 \log_2(1/\delta) \rfloor. \quad (\text{IV.40.8})$$

Proof. Make the same shear $z = x + w$ used in the repeated-variable argument. For a fixed w , combine the seven factors into

$$g_w(x) = b_1(x)b_3(x+w), \quad h_w(y) = b_2(y)b_7(y+w), \quad k_w(t) = b_4(t)b_6(t+w),$$

leaving a scalar factor $b_5(w)$. The tensor becomes

$$A_w(x, y) = T(x, x, y) + T(x, y, w).$$

It is bilinear, and its alternating part is Ω_T , independent of w . Its radical index parameter is therefore $d_{A_w} = 2^m$. Theorem IV.40.1 bounds each inner triangle correlation by $2^{-m/2}$. Averaging over w proves (IV.40.7). It follows that $m \leq 2 \log_2(1/\delta)$. An isotropic subspace of codimension m exists and has an integrable restriction, by the algebraic part of Theorem IV.39.4. $\square$

[write] **7 October 2026; sources 40–46.** The same codimension $\lfloor 2 \log_2(1/\delta) \rfloor$ is source 46’s Theorem IV.28.1, by a different route. Tidor’s Proposition 4.6 (re-read at the write in the arXiv HTML of version 2) gives codimension at most $8 \log_2(1/\delta)$ for symmetric trilinear forms over $\mathbb{F}_2$ under a seven-function correlation hypothesis; the hypotheses were compared only as displayed there, as both sources compare them.

Tidor’s Proposition 4.6 [150] has the same symmetric tensor and seven-function hypothesis, with codimension at most $8 \log_2(1/\delta)$. Here the fourth-power estimate controls the alternating rank more tightly, and an isotropic subspace replaces the full radical. The combined result is (IV.40.8). Optimality of the coefficient two for this seven-function problem is not established. In contrast, coefficient one in the one-function energy problem (IV.39.9) is proved optimal.

IV.41 The quartic frontier: an exact small quotient and a global reduction

Source 43, delivered Section 9.

IV.41.1 A useful nonintegrable family in every degree

For independent $u, v \in V^*$ and $d \geq 3$, put

$$C_d(h_1, \dots, h_d) = \sum_{i=1}^d v(h_i) \prod_{j \neq i} u(h_j). \quad (\text{IV.41.1})$$

Proposition IV.41.1 (Explicit lower bounds). *The tensor C_d is symmetric and nonintegrable, and*

$$\Omega_z(a, b) = \left(\prod_{j=1}^{d-3} u(z_j) \right) (u(a)v(b) + v(a)u(b)), \quad (\text{IV.41.2})$$

$$\mathcal{E}_{C_d}(1) = 1 - \frac{d+1}{2^d}. \quad (\text{IV.41.3})$$

Consequently the dimension-independent symmetric nonintegrable extremum satisfies

$$1 - \frac{d+1}{2^d} \leq \sup_{\substack{n \geq 2, \, T \in \text{SymMult}^d(\mathbb{F}_2^n) \setminus \text{nCSM}^d(\mathbb{F}_2^n) \\ \|f\|_\infty \leq 1}} \mathcal{E}_T(f) \leq 1 - 2^{2-d}. \quad (\text{IV.41.4})$$

Proof. In the two repeated evaluations, the terms where v is placed in one of the two repeated positions cancel in pairs. The terms where v is placed in a frozen position cancel between the two evaluations. The two remaining terms give (IV.41.2), which is nonzero for appropriate arguments.

For the constant function, the selected frequency is zero precisely when both coefficients in

$$y \mapsto v(y) \prod_{j=1}^{d-1} u(h_j) + u(y) \sum_{i=1}^{d-1} v(h_i) \prod_{j \neq i} u(h_j)$$

are zero. Put $k = d - 1$. If at least two of the $u(h_j)$ are zero, both coefficients vanish. If exactly one is zero, the remaining condition is one independent unbiased bit, so half the tuples work. If none is zero, the coefficient of $v(y)$ is one and the frequency cannot vanish. Therefore the success probability is

$$1 - 2^{-k} - \frac{1}{2}k2^{-k} = 1 - (d+1)2^{-d}.$$

The upper bound in (IV.41.4) is Corollary IV.38.3. $\square$

For $d = 3$ the bounds coincide at $1/2$. For $d = 4$ they give $11/16 \leq M \leq 3/4$. The next results explain why the fourth-degree problem is concentrated in one specific tensor rather than an unrestricted search over quartic coefficients.

IV.41.2 High-energy quartics have one rank-two defect

For a symmetric quartic, $z \mapsto \Omega_z$ is a linear map $V \rightarrow \text{Alt}^2(V)$. Denote its image dimension by r .

Theorem IV.41.2 (Quartic defect dichotomy and normal form). *Let $T : V^4 \rightarrow \mathbb{F}_2$ be symmetric and nonintegrable. If $\mathcal{E}_T(f) > 5/8$, then there are independent $u, v \in V^*$ and $S \in \text{nCSM}^4(V)$ such that*

$$T = S + C_4, \tag{IV.41.5}$$

where C_4 is (IV.41.1) for these two forms. More precisely, failure of this defect normal form implies the uniform bound $\mathcal{E}_T(f) \leq 5/8$.

Proof. Since T is nonintegrable, $r \geq 1$. If $r \geq 2$, the zero fiber of $z \mapsto \Omega_z$ has measure $2^{-r} \leq 1/4$. Every nonzero value has rank at least two. Hence the profile theorem gives

$$\mathcal{E}_T(f) \leq 2^{-r} + (1 - 2^{-r})/2 \leq 5/8.$$

Thus energy greater than $5/8$ forces $r = 1$. Write $\Omega_z = \ell(z)\Omega_0$ for nonzero ℓ and Ω_0 . If $\text{rank } \Omega_0 \geq 4$, then $\mathcal{E}_T(f) \leq (1 + 1/4)/2 = 5/8$. We must therefore have $\text{rank } \Omega_0 = 2$, so

$$\Omega_0(a, b) = u_0(a)v_0(b) + v_0(a)u_0(b)$$

for independent u_0, v_0 .

Symmetry of the original quartic imposes the identity

$$\Omega_c(a, b) + \Omega_b(a, c) + \Omega_a(b, c) = 0. \tag{IV.41.6}$$

Indeed, expansion gives six quartic terms, paired equally by symmetry. If ℓ were not in the span of u_0, v_0 , choose $c \in \ker u_0 \cap \ker v_0$ with $\ell(c) = 1$, and choose a, b with $\Omega_0(a, b) = 1$. Then the first term of (IV.41.6) would be one and the others zero, a contradiction. Thus $\ell \in \text{span}(u_0, v_0)$.

Set $u = \ell$. Because the scalar field is $\mathbb{F}_2$, one can choose v in the same two-dimensional span so that $\Omega_0 = u \otimes v + v \otimes u$. Formula (IV.41.2) for $d = 4$ shows that C_4 has exactly the same defect map as T . Their difference is symmetric and has zero repeated-variable defect, and hence lies in $\text{nCSM}^4(V)$. $\square$

The identity (IV.41.6) is important: a generic linear family of rank-two alternating forms need not be the defect of a symmetric quartic. Ignoring this compatibility would leave an unnecessarily large search space.

IV.41.3 The exact two-dimensional quartic extremum

Theorem IV.41.3 (The quartic value $11/16$ on $\mathbb{F}_2^2$). *For the tensor C_4 on $V = \mathbb{F}_2^2$ with $u(x) = x_1$ and $v(x) = x_2$,*

$$\max_{\|f\|_\infty \leq 1} \mathcal{E}_{C_4}(f) = \frac{11}{16}. \quad (\text{IV.41.7})$$

The same maximum holds for every symmetric nonintegrable quartic on $\mathbb{F}_2^2$, after an integrable phase twist and a linear change of coordinates.

Proof. Label the four points $0, e_1, e_2, e_1 + e_2$ by $0, 1, 2, 3$, and put $z_i = f(i)^4$. We first establish the exact polynomial identity

$$\begin{aligned} 1024\mathcal{E}_{C_4}(f) &= \left| z_0^2 + z_1^2 + z_2^2 + z_3^2 \right|^2 + 28 \sum_{i < j} |z_i z_j|^2 + 480 \prod_{i=0}^3 |z_i| \\ &\quad + 24 \operatorname{Re}(z_0 z_1 \overline{z_2 z_3} + z_0 z_3 \overline{z_1 z_2}) - 8 \operatorname{Re}(z_0 z_2 \overline{z_1 z_3}). \end{aligned} \quad (\text{IV.41.8})$$

A short cube census proves the formula without assumptions on the amplitudes. Here are all its cases.

If the four directions are all zero, their four possible base points give the diagonal terms of $|\sum_i f(i)^8|^2$. If the directions are all the same nonzero vector, their base points give its off-diagonal terms. All these tensor signs are positive. For a fixed nonzero vector, the fourteen remaining tuples in $\{0, h\}^4$ use both 0 and h . Each vertex of a two-point coset then occurs four times with each conjugation parity. The two possible base points per coset give the coefficient 28 of each $|f(i)f(j)|^8$.

It remains to consider direction tuples spanning $\mathbb{F}_2^2$. The cube map $\varepsilon \mapsto \sum_j \varepsilon_j h_j$ has four-point fibers. Unless the parity $\sum_j \varepsilon_j$ factors through that map, every fiber contains two even and two odd vertices, giving the common monomial $\prod_i |f(i)|^4$. Parity factors precisely when all four directions lie in a single two-element affine hyperplane not containing zero. They must use both of its points. For a pair of distinct nonzero directions a, b , the fourteen such tuples have multiplicity patterns $3 + 1$, $2 + 2$, and $1 + 3$. Their signed sum is

$$4(-1)^{C_4(a,a,a,b)} + 6(-1)^{C_4(a,a,b,b)} + 4(-1)^{C_4(a,b,b,b)}.$$

For the pairs (e_1, e_2) , $(e_1, e_1 + e_2)$, $(e_2, e_1 + e_2)$, these sums are $6, -2, 6$, respectively. Each orientation of the associated two-pair partition has two base points, giving coefficients $12, -4, 12$ and their conjugates. These are the three real-part terms in (IV.41.8).

Finally, the sum of the tensor signs over all direction tuples is $256(11/16) = 176$ by Proposition IV.41.1. The rank-zero and rank-one tuples contribute $1 + 3(15) = 46$; the parity-factoring

rank-two tuples contribute $6 - 2 + 6 = 10$. The remaining signed sum is therefore $176 - 46 - 10 = 120$. Their four base points give the coefficient 480. This completes the derivation of (IV.41.8).

The map $f(i) \mapsto z_i = f(i)^4$ maps the closed unit disc onto itself. The right side of (IV.41.8) is subharmonic separately in each z_i : the first term is a squared modulus of a holomorphic polynomial, the next two terms are nonnegative multiples of $|z_i|^2$ and $|z_i|$, and the remaining terms are harmonic in that variable. The disc maximum principle, applied one coordinate at a time, shows that its maximum is attained when every $|z_i| = 1$. This use of the maximum principle can also be seen directly from the Poisson mean inequality for the squared modulus and the nonnegative radial terms.

Write $z_i = e^{i\theta_i}$, and put

$$c = \cos(\theta_0 - \theta_2), \quad r = \cos(\theta_1 - \theta_3), \quad q = \cos(\theta_0 + \theta_2 - \theta_1 - \theta_3).$$

The second and third terms of (IV.41.8) contribute $168 + 480 = 648$. The remaining expression is

$$4(c^2 + r^2) + 48cr + 8q(cr - 1). \quad (\text{IV.41.9})$$

Since $cr \leq 1$ and $-1 \leq q \leq 1$,

$$4(c^2 + r^2) + 48cr + 8q(cr - 1) \leq 4(c^2 + r^2) + 40cr + 8 \leq 56.$$

Thus $1024\mathcal{E}_{C_4}(f) \leq 704$, or $\mathcal{E}_{C_4}(f) \leq 11/16$. The constant function attains this value.

On a two-dimensional space, $\text{Alt}^2(V)$ is one-dimensional. The nonzero defect map of any nonintegrable symmetric quartic therefore has image dimension one and nonzero rank two. The algebraic part of Theorem IV.41.2 applies without an energy assumption. It gives the stated normal form; gauge invariance and a linear change of variables preserve the maximum energy. $\square$

IV.41.4 The remaining higher-dimensional question is one fixed tensor

Let $C_{4,n}$ denote the pullback of C_4 to $\mathbb{F}_2^n = \mathbb{F}_2^2 \oplus \mathbb{F}_2^{n-2}$, and define

$$M_{4,n} = \max_{\substack{T \in \text{SymMult}^4(\mathbb{F}_2^n) \setminus \text{nCSM}^4(\mathbb{F}_2^n) \\ \|f\|_\infty \leq 1}} \mathcal{E}_T(f).$$

Corollary IV.41.4 (Exact reduction of the quartic optimization). *For every $n \geq 2$,*

$$M_{4,n} = \mathcal{M}(C_{4,n}), \quad \frac{11}{16} \leq M_{4,n} \leq \frac{3}{4}, \quad M_{4,2} = \frac{11}{16}. \quad (\text{IV.41.10})$$

The sequence $M_{4,n}$ is nondecreasing. Thus the full quartic extremal problem reduces to asking whether adjoining directions on which C_4 is zero increases its optimal tested energy.

Proof. The explicit tensor has energy $11/16 > 5/8$ at the constant function. Any tensor whose defect does not have the normal form has energy at most $5/8$, so it cannot determine the maximum. Every remaining tensor is, after gauge and coordinates, exactly $C_{4,n}$. This proves the equality of optimization problems. The lower and upper bounds follow from the explicit example and the shear profile. An optimizer on dimension n may be pulled back while ignoring an additional coordinate, proving monotonicity. $\square$

We do not infer $M_{4,n} = M_{4,2}$ from the fact that the tensor uses only two coordinates. The optimizing function may depend on all the coordinates. Ruling out such a gain is the precise additional theorem that would be needed. The cubic tensorization theorem does not settle this quartic question.

IV.42 Source 47's chapter: exact Boolean obstruction energies [write]

Added in the ninth write, 7 October 2026.

The next eight sections are source 47's Sections 2–9. *Logically they continue source 43's part of the binary chapter* (Section IV.25), whose quartic question they answer; they are printed here because no number printed earlier may move. Source 47's introduction is Section F.52.

What is proved. Over $V = \mathbb{F}_2^n$, for the canonical symmetric tensor $C_d = \sum_i v(h_i) \prod_{j \neq i} u(h_j)$ with independent u, v , the maximum selected energy is exactly $1 - (d+1)/2^d$ in every degree $d \geq 3$ and every dimension $n \geq 2$, attained by the constant function; optimizers may use all coordinates (Theorem F.52.1). The key is that every even sum of derivative directions is an exact period, which forbids frequencies (Section IV.45). For all symmetric nonintegrable tensors the maximum energy lies between $1 - (d+1)/2^d$ and $1 - \min\{d+1, 7\}/2^d$, so the thresholds are exactly $1/2, 11/16, 13/16, 57/64$ in degrees three to six (Theorem F.52.2), through a pure-defect classification and a sharp $7/32$ support bound for non-pure symmetric trilinear maps (Theorem IV.48.7); above the threshold the tensor integrates and phase removal is lossless. A nonnegative deficit recursion and robust certification for corrupted selectors (Propositions IV.50.1, IV.50.2). At $d = 4$ this is source 43's quartic question, answered (item IV.10); the septic value is open ($120/128 \leq M_{7,n}^{\text{sym}} \leq 121/128$).

Relation to the formal statements (checked at HEAD 06b30ca52). `lemma_17_1` and `proposition_17_2` (Sections17_18.lean, lines 38 and 48) have the exact companions `lemma_17_1_holds` (Proofs17Phase.lean, line 374) and `proposition_17_2_holds` (Proofs17PhaseRemoval.lean, line 373); they are stated on $\mathbb{Z}/N\mathbb{Z}$ with N prime, so characteristic two lies outside them, as source 47 says. Nothing in this chapter is formalized.

IV.43 Finite calculus and nonclassical integration

Source 47, delivered Section 2.

IV.43.1 The cube identity and exact slicing

Lemma IV.43.1 (Selected cube identity). *For every complex f and every multilinear T ,*

$$\mathcal{E}_T(f) = \mathbb{E}_{x, h_1, \dots, h_d} \partial_{h_1} \cdots \partial_{h_d} f(x) (-1)^{T(h_1, \dots, h_d)}. \quad (\text{IV.43.1})$$

For symmetric T and $d \geq 3$, if $T_z(h_1, \dots, h_{d-1}) = T(h_1, \dots, h_{d-1}, z)$, then

$$\mathcal{E}_T(f) = \mathbb{E}_z \mathcal{E}_{T_z}(\partial_z f). \quad (\text{IV.43.2})$$

Proof. Fix $\mathbf{h} = (h_1, \dots, h_{d-1})$ and write $g = \partial_{h_1} \cdots \partial_{h_{d-1}} f$. Expanding the squared Fourier coefficient and substituting one base point as the other plus y gives

$$\left| \mathbb{E}_x g(x) (-1)^{T(\mathbf{h}, x)} \right|^2 = \mathbb{E}_{x, y} g(x+y) \overline{g(x)} (-1)^{T(\mathbf{h}, y)}.$$

Multiplicative derivatives commute. Averaging proves (IV.43.1). In the symmetric case freeze the last direction as z , commute its derivative to the function, and apply the cube identity in degree $d-1$. This proves (IV.43.2). $\square$

The cube average is real and nonnegative because its original expression is an average of squares. Define

$$\|f\|_{U^d}^{2^d} = \mathcal{E}_0(f).$$

For the purposes of this article, only this identity is needed, not the full norm theory.

Lemma IV.43.2 (Gauge invariance). *If $S \in \mathcal{I}_d(V)$ and P satisfies (F.52.2) with S in place of T , then*

$$\mathcal{E}_{T+S}(fe(P)) = \mathcal{E}_T(f). \quad (\text{IV.43.3})$$

In particular $\mathcal{M}(T + S) = \mathcal{M}(T)$.

Proof. The d -fold derivative of $e(P)$ is $(-1)^S$. In the cube identity the extra two factors $(-1)^S$ cancel. Multiplication by $e(P)$ is a bijection of the unit polydisc of functions, proving the assertion about maxima. $\square$

These identities remain valid when f vanishes; there is no division by f anywhere in the argument.

IV.43.2 The precise integration criterion

For $d \geq 3$, define the repeated-variable defect of symmetric T by

$$\Phi_T(z_1, \dots, z_{d-3})(a, b) = T(a, a, b, z_1, \dots, z_{d-3}) + T(a, b, b, z_1, \dots, z_{d-3}). \quad (\text{IV.43.4})$$

For $d = 3$ this is simply an alternating bilinear form $\Omega_T(a, b)$.

Theorem IV.43.3 (Boolean nonclassical integration). *A symmetric d -linear tensor, $d \geq 3$, is in $\mathcal{I}_d(V)$ if and only if $\Phi_T = 0$. Every symmetric bilinear tensor is integrable. In coordinates $x_1, \dots, x_n$, a primitive can be chosen as*

$$P_T(x) = \sum_{\substack{\emptyset \neq A \subseteq [n] \\ |A| \leq d}} \frac{c_A \prod_{i \in A} |x_i|}{2^{d-|A|+1}} \pmod{1}, \quad (\text{IV.43.5})$$

where c_A is the value of T on any length- d basis tuple whose set of indices is exactly A , and $|x_i| \in \{0, 1\}$.

This is the Boolean specialization of the known nonclassical integration criterion [150, Definition 3.1 and Proposition 3.5]; we give a direct coordinate proof rather than invoke that result as a black box.

Proof. On functions taking values in $\mathbb{R}/\mathbb{Z}$, translation by a has square the identity. Consequently

$$\Delta_a^2 = -2\Delta_a, \quad \Delta_a^2 \Delta_b = \Delta_a \Delta_b^2. \quad (\text{IV.43.6})$$

Commuting additive differences proves symmetry and the required repeated-variable identity for every top derivative. This proves necessity.

Conversely, assume $\Phi_T = 0$. Symmetry makes a basis coefficient depend only on its multiplicities. When two indices i, j occur and i occurs at least twice, the identity $T(e_i, e_i, e_j, \mathbf{z}) = T(e_i, e_j, e_j, \mathbf{z})$ moves one occurrence from i to j without changing the coefficient. All positive multiplicity lists of total d on a fixed support are connected by these moves. Thus c_A is well defined.

For a set A of size a , consider $P_A(x) = 2^{-(d-a+1)} \prod_{i \in A} |x_i|$. Repeated coordinate differences satisfy $\Delta_{e_i}^t = (-2)^{t-1} \Delta_{e_i}$ for $t \geq 1$. Moreover

$$\Delta_h = \prod_{i: h_i=1} (I + \Delta_{e_i}) - I. \quad (\text{IV.43.7})$$

A term in the expansion of $d+1$ differences contains at least $d+1$ coordinate differences. If a coordinate outside A occurs, the term is zero. Otherwise, if $r \leq a$ distinct coordinates occur among $t \geq d+1$ differences, its numerator is divisible by 2^{t-r} , hence by 2^{d-a+1} . It vanishes modulo one. Therefore P_A has degree at most d .

Apply exactly d coordinate differences. A support properly contained in A again produces an integer, and a coordinate outside A produces zero. On a tuple of support exactly A , the result is $\pm 1/2 = 1/2 \pmod{1}$, independent of the base point. Thus the top derivative of (IV.43.5) has the required basis coefficients. The top derivative of a degree-at-most- d polynomial is multiadditive: $\Delta_{a+b} = \Delta_a + \Delta_b + \Delta_a \Delta_b$, and the error in additivity is a $(d+1)$ -fold derivative. It is independent of the base point for the same reason. Agreement on basis tuples proves the result. For $d=2$ the same construction works without a multiplicity-transfer condition. $\square$

Corollary IV.43.4 (Pure powers integrate). *For $u \in V^*$ and $m \geq 2$, set $U_m(h_1, \dots, h_m) = \prod_i u(h_i)$. Then*

$$P_{u,m}(x) = \frac{|u(x)|}{2^m} \pmod{1} \quad (\text{IV.43.8})$$

is a primitive of U_m .

Proof. Apply the one-coordinate case of (IV.43.5) and pull back along u . $\square$

IV.44 The spectral obstruction supplied by the defect

Source 47, delivered Section 3.

This section reconstructs the shear method of the preceding draft [205]. It supplies the cubic base case and a profile estimate used later. The new improvements will come from preserving derivative constraints and from stronger information about the support of the defect.

IV.44.1 A finite projective orthogonality bound

Let $B : V \times V \rightarrow \mathbb{F}_2$ be bilinear, and put

$$\Omega_B(a, b) = B(a, b) + B(b, a), \quad \text{rank } \Omega_B = 2r.$$

An alternating bilinear form has even rank over $\mathbb{F}_2$, as is seen by successively splitting off a pair a, b with $\Omega_B(a, b) = 1$. Define

$$\mathcal{Q}_B(g, h) = \mathbb{E}_a \left| \mathbb{E}_x g(x+a) \overline{h(x)} (-1)^{B(a,x)} \right|^2, \quad \mathcal{Q}_B(g) = \mathcal{Q}_B(g, g). \quad (\text{IV.44.1})$$

Lemma IV.44.1 (Projective orthogonality). *For arbitrary complex functions g, h ,*

$$\mathcal{Q}_B(g, h) \leq 2^{-r} \|g\|_2^2 \|h\|_2^2. \quad (\text{IV.44.2})$$

Also

$$\mathbb{E}_{a,b,x} \partial_a \partial_b g(x) (-1)^{B(a,b)} = \mathcal{Q}_B(g). \quad (\text{IV.44.3})$$

Proof. Use the normalized Hilbert space inner product $\langle g, h \rangle = \mathbb{E}_x g(x) \overline{h(x)}$, linear in the first variable. The operators

$$(U_a g)(x) = (-1)^{B(a,x)} g(x+a)$$

are unitary and satisfy

$$U_a U_b = (-1)^{B(b,a)} U_{a+b}, \quad U_a U_b = (-1)^{\Omega_B(a,b)} U_b U_a.$$

The positive operator

$$M_g v = \mathbb{E}_a \langle v, U_a g \rangle U_a g$$

has trace $\|g\|_2^2$ and commutes with every U_b : conjugation by U_b permutes its rank-one summands, and the scalar signs disappear.

Choose a subspace H containing the radical of Ω_B whose image in the symplectic quotient is Lagrangian. Then $H = H^\perp$ and $[V : H] = 2^r$. Every nonzero common invariant subspace of the U_a has dimension at least 2^r . Indeed the commuting unitaries U_h , $h \in H$, have a common eigenvector in that subspace. Its translates by representatives of V/H have different joint eigenvalues, since their ratios are the characters $h \mapsto (-1)^{\Omega_B(h,a)}$. They are pairwise orthogonal.

Every nonzero eigenspace of M_g is such an invariant subspace. Thus each positive eigenvalue has multiplicity at least 2^r , and $\|M_g\|_{\text{op}} \leq 2^{-r} \text{tr } M_g$. Since $\mathcal{Q}_B(g, h) = \langle M_g h, h \rangle$, this proves (IV.44.2). Finally, expand the square in (IV.44.1) and change one base point to the other plus b . The resulting average is the complex conjugate of the left side of (IV.44.3); it is real because it equals $\mathcal{Q}_B(g)$. This proves the identity. $\square$

IV.44.2 The exact shear and its rank profile

Lemma IV.44.2 (Bilinearity of the defect). *For symmetric T , the map $\Phi_T : V^{d-3} \rightarrow \text{Alt}^2(V)$ is symmetric multilinear in its frozen variables. For each frozen tuple its value is an alternating bilinear form in a, b .*

Proof. In $T(a + a', a + a', b, z)$ the two cross terms cancel by symmetry. This proves additivity in the repeated variable, and the other variables are already linear. The same argument applies to the repeated b . The value at $a = b$ is zero. Linearity and symmetry in the frozen variables are inherited directly from T . $\square$

Theorem IV.44.3 (Shear profile). *Let T be symmetric of degree $d \geq 3$. With $\tau_w f(x) = f(x + w)$, put*

$$g_{z,w} = \partial_{z_1} \cdots \partial_{z_{d-3}} (f \tau_w f), \quad B_{z,w}(a, b) = T(a, a, b, z) + T(a, b, z, w).$$

Then

$$\mathcal{E}_T(f) = \mathbb{E}_{z,w} \mathcal{Q}_{B_{z,w}}(g_{z,w}). \quad (\text{IV.44.4})$$

Consequently, for $\|f\|_\infty \leq 1$,

$$\mathcal{E}_T(f) \leq \mathbb{E}_z 2^{-\text{rank } \Phi_T(z)/2}. \quad (\text{IV.44.5})$$

For $d = 3$ the right side is $2^{-\text{rank } \Omega_T/2}$.

Proof. The pointwise identity

$$\partial_a \partial_{a+w} f = \overline{\partial_a (f \tau_w f)} \quad (\text{IV.44.6})$$

follows by multiplying the four factors. In the selected cube identity, write the directions as a, b, z, c and substitute $c = a + w$. Multilinearity gives the form $B_{z,w}$ above. It is bilinear by the same cancellation as in Lemma IV.44.2, and its alternating part is $\Phi_T(z)$, since the term containing w is symmetric in a, b . Apply the remaining derivatives to (IV.44.6) and then use (IV.44.3). Its reality removes the conjugation and proves (IV.44.4). Orthogonality bounds each summand by $2^{-\text{rank } \Phi_T(z)/2} \|g_{z,w}\|_2^4$. The last norm is at most one. $\square$

Corollary IV.44.4 (Support-to-energy inequality). *For symmetric nonintegrable T of degree $d \geq 4$,*

$$\mathcal{E}_T(f) \leq 1 - \frac{1}{2} \mathbb{P}_z[\Phi_T(z) \neq 0]. \quad (\text{IV.44.7})$$

Proof. Every nonzero alternating form has rank at least two and thus contributes at most $1/2$ to (IV.44.5). $\square$

IV.45 Derivative periods and forbidden frequencies

Source 47, delivered Section 4.

IV.45.1 The even span is an exact period space

For $\mathbf{h} = (h_1, \dots, h_r)$ define

$$H_{\text{ev}}(\mathbf{h}) = \left\{ \sum_{i=1}^r \varepsilon_i h_i : \varepsilon_i \in \mathbb{F}_2, \sum_{i=1}^r \varepsilon_i = 0 \right\}. \quad (\text{IV.45.1})$$

It is the span of $h_1 + h_i$ for $2 \leq i \leq r$.

Lemma IV.45.1 (Exact period constraint). *For every complex f , including functions with zeros, and $r \geq 1$, let $F_{\mathbf{h}} = \partial_{h_1} \cdots \partial_{h_r} f$. Then*

$$F_{\mathbf{h}}(x + h_i) = \overline{F_{\mathbf{h}}(x)} \quad (1 \leq i \leq r), \quad F_{\mathbf{h}}(x + w) = F_{\mathbf{h}}(x) \quad (w \in H_{\text{ev}}(\mathbf{h})). \quad (\text{IV.45.2})$$

In particular, for a frequency $\xi \in V^$,*

$$\widehat{F_{\mathbf{h}}}(\xi) \neq 0 \implies \xi(h_1) = \xi(h_2) = \cdots = \xi(h_r). \quad (\text{IV.45.3})$$

Proof. The product formula for $F_{\mathbf{h}}(x)$ has one factor at every labelled cube vertex, conjugated according to the parity of $r - |\varepsilon|$. Translation by h_i toggles the i th vertex coordinate and reverses conjugation parity. It therefore conjugates the whole product. Two such translations preserve the product, proving invariance under the even span. This is an identity of products, even when some labelled vertices coincide.

If $w \in H_{\text{ev}}(\mathbf{h})$ and $\xi(w) = 1$, translation by w changes the sign of the Fourier character while leaving the function unchanged. The coefficient equals its negative and is zero. Annihilating the generators $h_1 + h_i$ is exactly the condition in (IV.45.3). $\square$

For two directions the identity has the particularly transparent form

$$\partial_a \partial_z f(x) = f(x) f(x + a + z) \overline{f(x + a) f(x + z)} = \partial_a \partial_z f(x + a + z).$$

No approximation, exceptional set, or unit-modulus hypothesis is involved.

Corollary IV.45.2 (A support bound for arbitrary selectors). *For an arbitrary selector $\sigma : V^r \rightarrow V^*$ and $\|f\|_\infty \leq 1$,*

$$\begin{aligned} \mathbb{E}_{\mathbf{h}} \left| \widehat{\partial_{\mathbf{h}} f}(\sigma(\mathbf{h})) \right|^2 \\ \leq \mathbb{E}_{\mathbf{h}} \mathbf{1}_{\{\sigma(\mathbf{h})|_{H_{\text{ev}}(\mathbf{h})} = 0\}} \|\partial_{\mathbf{h}} f\|_2^2 \leq \mathbb{P}_{\mathbf{h}}[\sigma(\mathbf{h})|_{H_{\text{ev}}(\mathbf{h})} = 0]. \end{aligned} \quad (\text{IV.45.4})$$

Proof. The inadmissible coefficients vanish by Lemma IV.45.1. For each admissible coefficient apply Cauchy–Schwarz, and then use $|\partial_{\mathbf{h}} f| \leq 1$. $\square$

The bound can be weaker than the shear profile if applied indiscriminately to the original tensor. Its decisive use here is on a slice that is known to be a derivative of the *same* original function.

IV.45.2 The sharp derivative-constrained rank-one estimate

Theorem IV.45.3 (A rank-one tensor tested on a fixed derivative). *Let $u \in V^*$ be nonzero, $m \geq 2$, and $z \in \ker u$. For every $\|f\|_\infty \leq 1$,*

$$\boxed{\mathcal{E}_{U_m}(\partial_z f) \leq 1 - 2^{1-m}, \quad U_m(h_1, \dots, h_m) = \prod_i u(h_i).} \quad (\text{IV.45.5})$$

The bound is attained by $f = 1$, for every such z and in every dimension.

Proof. In the selected-energy definition for U_m , the frequency is u when $u(h_i) = 1$ for all $1 \leq i \leq m-1$, and is zero otherwise. On the former set the function being transformed is

$$\partial_{h_1} \cdots \partial_{h_{m-1}} \partial_z f.$$

It is invariant under $h_1 + z$, while $u(h_1 + z) = 1$. Its selected coefficient therefore vanishes. The excluded set has probability $2^{-(m-1)}$; all remaining squared coefficients are at most one. This proves the bound. For $f = 1$ all derivatives are one, so every zero-frequency coefficient is one and every nonzero-frequency coefficient is zero. Equality follows. $\square$

Remark IV.45.4 (Why unrestricted slice optimization loses the answer). The tensor U_m itself is integrable, so $\mathcal{M}(U_m) = 1$, attained by $e(P_{u,m})$. Theorem IV.45.3 says that this value cannot be attained inside the constrained family $\{\partial_z f : |f| \leq 1\}$ when $u(z) = 0$. Replacing a derivative slice by an arbitrary bounded function erases a loss of exactly 2^{1-m} . For the quartic problem, that loss occurs on one quarter of the slices and supplies the missing $1/16$.

IV.46 Exact canonical energies in every degree

Source 47, delivered Section 5.

IV.46.1 The defect and the lower bound

Lemma IV.46.1 (Canonical defect and constant energy). *For $d \geq 3$,*

$$\Phi_{C_d}(\mathbf{z})(a, b) = \left(\prod_{j=1}^{d-3} u(z_j) \right) (u(a)v(b) + v(a)u(b)), \quad (\text{IV.46.1})$$

with the empty product equal to one. In particular C_d is not integrable. Moreover

$$\mathcal{E}_{C_d}(1) = 1 - \frac{d+1}{2^d} = b_d. \quad (\text{IV.46.2})$$

Proof. In $C_d(a, a, b, \mathbf{z})$, the two terms placing v in a repeated a slot cancel. The terms placing v in a frozen slot cancel against their counterparts in $C_d(a, b, b, \mathbf{z})$. The two remaining terms give (IV.46.1). Independence of u, v makes the defect nonzero.

For the constant function, the energy is the probability that $C_d(h_1, \dots, h_{d-1}, \cdot)$ is the zero functional. Its coefficients of v and u are respectively

$$A = \prod_{j=1}^{d-1} u(h_j), \quad B = \sum_{i=1}^{d-1} v(h_i) \prod_{j \neq i} u(h_j).$$

If at least two of the $u(h_j)$ vanish, both coefficients vanish. If exactly one vanishes, $A = 0$ and B is the corresponding independent unbiased v bit. If none vanishes, $A = 1$. With $k = d - 1$, the failure probability is therefore $2^{-k} + \frac{1}{2}k2^{-k} = (d+1)2^{-d}$. $\square$

This lower-bound computation and the canonical defect were already present in the preceding draft [205]. The upper bound that follows is the additional all-dimensional result.

IV.46.2 The inductive upper bound

Proof of Theorem F.52.1. For $d = 3$, the alternating defect is $u \otimes v + v \otimes u$, of rank two. The shear profile gives $\mathcal{E}_{C_3}(f) \leq 1/2 = b_3$. The lower bound is Lemma IV.46.1.

Assume the result in degree $m = d - 1 \geq 3$. Exact slicing gives

$$\mathcal{E}_{C_d}(f) = \mathbb{E}_z \mathcal{E}_{u(z)C_m + v(z)U_m}(\partial_z f), \quad (\text{IV.46.3})$$

since each term of C_d either places v in the last slot or does not. The four pairs $(u(z), v(z))$ are equidistributed.

On the quarter $(0, 0)$ the tensor is zero and the slice energy is at most one. On the quarter $(0, 1)$ it is U_m , and Theorem IV.45.3 gives the sharper bound $1 - 2^{1-m}$ because $u(z) = 0$. On the half $u(z) = 1$, the tensor is either C_m or $C_m + U_m$. In the latter case remove the integrable tensor U_m using Corollary IV.43.4 and gauge invariance. The resulting function is still one-bounded, so the induction hypothesis gives b_m in both cases.

It follows that

$$\begin{aligned} \mathcal{E}_{C_d}(f) &\leq \frac{1}{4} + \frac{1}{4}(1 - 2^{2-d}) + \frac{1}{2}b_{d-1} \\ &= \frac{1}{4} + \frac{1}{4} - 2^{-d} + \frac{1}{2} \left(1 - \frac{d}{2^{d-1}}\right) = 1 - \frac{d+1}{2^d}. \end{aligned} \quad (\text{IV.46.4})$$

Together with the constant lower bound, this proves equality in every degree. Gauge invariance gives the assertion for $C_d + S$. $\square$

IV.46.3 The quartic stabilization question, explicitly

The preceding draft left open whether the pullback $C_{4,n}$ of its binary quartic to $\mathbb{F}_2^n$ could gain energy from extra coordinates in f . For $d = 4$, the preceding proof specializes to

$$\mathcal{E}_{C_{4,n}}(f) \leq \frac{1}{4} \cdot 1 + \frac{1}{4} \cdot \frac{3}{4} + \frac{1}{2} \cdot \frac{1}{2} = \frac{11}{16}. \quad (\text{IV.46.5})$$

The $3/4$ term is the constrained pure-cubic slice, not an unrestricted cubic maximum. Thus this argument supplies exactly the missing theorem; no assertion about the dependence of an optimizer on the inactive coordinates is needed. The same mechanism gives $13/16$ in degree five and $57/64$ in degree six for the canonical family before any classification of other tensors is considered.

IV.47 The geometry of low-complexity defects

Source 47, delivered Section 6.

IV.47.1 A compatibility identity

Lemma IV.47.1 (Bianchi-type identity). *For $d \geq 4$ and each fixed tuple $\mathbf{w} \in V^{d-4}$,*

$$\Phi_T(c, \mathbf{w})(a, b) + \Phi_T(b, \mathbf{w})(a, c) + \Phi_T(a, \mathbf{w})(b, c) = 0. \quad (\text{IV.47.1})$$

Proof. Expand the three defects. There are six values of T ; symmetry pairs them equally, so their sum is zero in $\mathbb{F}_2$. $\square$

Lemma IV.47.2 (A one-form compatibility test). *Let $u \in V^*$ be nonzero and $\Omega \in \text{Alt}^2(V)$. If*

$$u(c)\Omega(a, b) + u(b)\Omega(a, c) + u(a)\Omega(b, c) = 0 \quad \text{for all } a, b, c, \quad (\text{IV.47.2})$$

then $\Omega = u \otimes v + v \otimes u$ for some $v \in V^$. If $\Omega \neq 0$, u, v are independent and $\text{rank } \Omega = 2$.*

Proof. Choose c_0 with $u(c_0) = 1$ and set $v(x) = \Omega(x, c_0)$. Substitute $c = c_0$ into (IV.47.2). It gives exactly $\Omega(a, b) = u(a)v(b) + v(a)u(b)$. A nonzero form of this shape has rank two. $\square$

Thus a rank-four alternating form cannot occur with just one nonzero linear frozen factor. The compatibility identity is stronger than merely knowing that each defect value is alternating.

IV.47.2 Pure symmetric maps and the canonical class

Definition IV.47.3 (Pure symmetric multilinear map). A nonzero symmetric s -linear map $F : V^s \rightarrow W$ between finite-dimensional Boolean vector spaces is *pure* if

$$F(z_1, \dots, z_s) = A \prod_{i=1}^s u(z_i) \quad (\text{IV.47.3})$$

for some nonzero $A \in W$ and nonzero $u \in V^*$. For $s = 1$, this means that F has rank one. A nonzero map that is not pure is called non-pure.

Theorem IV.47.4 (Pure defects are exactly canonical obstructions). *Let $d \geq 4$ and $T \in \text{SymMult}^d(V)$ be nonintegrable. The following are equivalent:*

- (i) $\Phi_T : V^{d-3} \rightarrow \text{Alt}^2(V)$ is pure;
- (ii) $T = C_d^{u,v} + S$ for independent $u, v \in V^*$ and $S \in \mathcal{I}_d(V)$;
- (iii) the first-slot radical of Φ_T has codimension one.

Proof. Put $s = d - 3$. If $\Phi_T(\mathbf{z}) = \Omega \prod_i u(z_i)$, choose all the $s - 1$ remaining frozen directions to have u value one in (IV.47.1). Lemma IV.47.2 gives $\Omega = u \otimes v + v \otimes u$ with u, v independent. By (IV.46.1), $T - C_d^{u,v}$ has zero defect and is integrable. This proves (i) $\Rightarrow$ (ii); the reverse implication is the same explicit defect formula.

For any symmetric multilinear map, its first-slot radical is a radical in every slot. If that radical has codimension one, the map descends to an s -linear map on a one-dimensional quotient and necessarily has the form (IV.47.3). Conversely a nonzero pure map has radical $\ker u$, proving the equivalence with (iii). $\square$

In particular, being pure is much stronger than having a one-dimensional *image* when $s > 1$. The latter case still admits a useful description.

Theorem IV.47.5 (One-dimensional defect images have two active coordinates). *Let $d \geq 4$ and suppose the linear span of the values of Φ_T is one-dimensional and nonzero. Then there is a surjection $q : V \rightarrow \mathbb{F}_2^2$, a symmetric d -linear T_0 on $\mathbb{F}_2^2$, and $S \in \mathcal{I}_d(V)$ such that*

$$T = S + q^*T_0. \quad (\text{IV.47.4})$$

The common nonzero alternating value of the defect has rank two.

Proof. Write $\Phi_T(\mathbf{z}) = F(\mathbf{z})\Omega_0$, where F is a nonzero scalar symmetric $(d - 3)$ -linear map. Some first-slot contraction $\ell(\mathbf{z}) = F(\mathbf{z}, \mathbf{w})$ is nonzero. By the Bianchi identity and Lemma IV.47.2, Ω_0 has rank two. Write $\Omega_0 = u \otimes v + v \otimes u$ with u, v independent, and put $R = \ker u \cap \ker v$.

Every first-slot contraction ℓ of F vanishes on R . Indeed, for $c \in R$ choose a, b with $\Omega_0(a, b) = 1$ in (IV.47.2); the identity gives $\ell(c) = 0$. Symmetry gives the same conclusion in every frozen slot. Thus the whole defect, including its alternating pair, factors through $q = (u, v) : V \rightarrow \mathbb{F}_2^2$.

Choose a linear section $j : \mathbb{F}_2^2 \rightarrow V$ and let $T_0 = j^*T$. The defect of q^*T_0 is the same as that of T , since qj is the identity and the defect factors through q . Their difference is integrable by Theorem IV.43.3. $\square$

This reduces the *tensor*, not the optimizing function. It must not be used to infer $\mathcal{M}(q^*T_0) = \mathcal{M}(T_0)$ for general T_0 . Theorem F.52.1 establishes that equality for the canonical family by an additional argument.

IV.47.3 A retained radical-sensitive profile

Let $F : V^s \rightarrow W$ be nonzero and symmetric, and let r be the codimension of its first-slot radical. A nonzero t -linear map over $\mathbb{F}_2$ is nonzero on at least a 2^{-t} fraction of its inputs: freeze all but one variable and induct. Consequently

$$\mathbb{P}[F \neq 0] \geq (1 - 2^{-r})2^{-(s-1)}. \quad (\text{IV.47.5})$$

Applying this to $F = \Phi_T$, $s = d - 3$, gives

$$\mathcal{E}_T(f) \leq 1 - 2^{3-d}(1 - 2^{-r}). \quad (\text{IV.47.6})$$

This bound can be stronger than our dimension-independent bound when r is large. For $r = 1$, the exact canonical result is stronger and should be used instead. The next section improves the small- r , non-pure case.

IV.48 Support bounds for non-pure symmetric multilinear maps

Source 47, delivered Section 7.

The following finite-field estimates are independent of Fourier analysis. They apply to vector-valued maps, which is essential when the target is $\text{Alt}^2(V)$.

IV.48.1 Two structural reductions

Lemma IV.48.1 (Distinct pure factors do not add to a pure tensor). *Let $m \geq 2$, $A, B \neq 0$ in W , and let $u, v \in V^*$ be distinct nonzero forms. Then $Au^{\otimes m} + Bv^{\otimes m}$ is neither zero nor pure.*

Proof. Over $\mathbb{F}_2$, u, v are independent. Choose a, b with $(u(a), v(a)) = (1, 0)$ and $(u(b), v(b)) = (0, 1)$. If the sum were $Dw^{\otimes m}$, evaluation on $(a, \dots, a)$ and $(b, \dots, b)$ would force $D = A = B \neq 0$ and $w(a) = w(b) = 1$. Evaluation on a tuple containing both a and b would then give zero on the left and D on the right. The same first evaluation rules out the zero tensor. $\square$

Lemma IV.48.2 (Non-purity survives a scalar projection). *If $s \geq 2$ and $F : V^s \rightarrow W$ is non-pure, some linear functional $\lambda \in W^*$ makes $\lambda \circ F$ non-pure.*

Proof. Suppose every scalar projection is zero or pure. The image of $W^* \rightarrow \text{SymMult}^s(V)$ is a linear subspace. It cannot contain two distinct nonzero pure scalar tensors, because their sum is non-pure by Lemma IV.48.1. Hence that image has dimension at most one. Since $F \neq 0$, it is generated by some $u^{\otimes s}$, and separation by linear functionals gives $F = Au^{\otimes s}$ for some $A \neq 0$. This contradicts non-purity. $\square$

Lemma IV.48.3 (Non-purity survives a contraction). *If $s \geq 3$ and $F : V^s \rightarrow W$ is symmetric and non-pure, there exists z such that $F_z = F(z, \cdot, \dots, \cdot)$ is non-pure.*

Proof. Suppose all nonzero slices are pure. Since $z \mapsto F_z$ is linear, Lemma IV.48.1 implies that their pure factors are all the same nonzero $u \in V^*$. Thus $F_z = A_z u^{\otimes(s-1)}$, where $z \mapsto A_z$ is linear. Choose w with $u(w) = 1$. Symmetry gives

$$A_z = F(z, w, \dots, w) = F(w, z, w, \dots, w) = u(z)A_w.$$

It follows that $F = A_w u^{\otimes s}$, contradicting non-purity. $\square$

Lemma IV.48.4 (Support lifting). *For a multilinear map $F : V^s \rightarrow W$ and any fixed z ,*

$$\mathbb{P}[F \neq 0] \geq \frac{1}{2} \mathbb{P}[F_z \neq 0]. \quad (\text{IV.48.1})$$

Proof. Fix the other inputs. If their slice has a nonzero value at z , the linear map in the first input is nonzero and hence nonzero on at least half the first inputs. Average this pointwise implication. $\square$

IV.48.2 The sharp bilinear and trilinear support bounds

Proposition IV.48.5 (Non-pure bilinear support). *A symmetric non-pure bilinear map $F : V^2 \rightarrow W$ is nonzero on at least a $3/8$ fraction of its inputs. The constant is sharp.*

Proof. The first-slot radical cannot have codimension one, since symmetry would then make F pure. Thus its codimension r is at least two, and (IV.47.5) gives $(1 - 2^{-r})/2 \geq 3/8$. The scalar form $u(a)u(b) + v(a)v(b)$ for independent u, v has exactly this support. $\square$

Lemma IV.48.6 (The binary trilinear census). *For a scalar symmetric trilinear form on $\mathbb{F}_2^2$, write*

$$A = F(e_1, e_1, e_1), \quad B = F(e_1, e_1, e_2), \quad C = F(e_1, e_2, e_2), \quad D = F(e_2, e_2, e_2).$$

If the form is nonzero and non-pure, its support has size at least 14 out of 64 input triples. The complete census is as follows.

Support size	Coefficient strings $ABCD$	Description
0	0000	zero
8	0001, 1000, 1111	pure
14	0111, 1001, 1110	non-pure
16	0010, 0011, 0100, 0101, 1010, 1100	non-pure
18	0110, 1011, 1101	non-pure

Proof. A triple with a zero input contributes zero. Write the three nonzero vectors as $a, b, c = a + b$. The three constant triples have values $A, D, A + B + C + D$. The six types with a repeated input have values $B, C, A + B, A + C, B + D, C + D$, each with multiplicity three. The all-distinct type has value $B + C$ and multiplicity six. Therefore the support size is

$$\begin{aligned} & [A] + [D] + [A + B + C + D] \\ & + 3([B] + [C] + [A + B] + [A + C] + [B + D] + [C + D]) + 6[B + C], \end{aligned} \quad (\text{IV.48.2})$$

where brackets denote the 0–1 representative of a Boolean value. Substitution of the sixteen coefficient strings gives the displayed table. The three nonzero pure forms are exactly $u^{\otimes 3}$ for the three nonzero linear forms u on $\mathbb{F}_2^2$. This is a complete finite arithmetic proof; the accompanying program independently checks every row. $\square$

Theorem IV.48.7 (Sharp non-pure trilinear support). *For every symmetric non-pure trilinear map $F : V^3 \rightarrow W$,*

$$\mathbb{P}[F \neq 0] \geq \frac{7}{32}. \quad (\text{IV.48.3})$$

The coefficient is sharp, even for scalar-valued maps on $\mathbb{F}_2^2$.

Proof. By Lemma IV.48.2, choose a non-pure scalar projection $G = \lambda \circ F$. Its support is contained in that of F , so it is enough to treat G . Let r be the codimension of its first-slot radical. For $r = 1$, symmetry would make it pure, which is excluded. For $r \geq 3$, (IV.47.5) yields

$$\mathbb{P}[G \neq 0] \geq (1 - 2^{-r}) \cdot \frac{1}{4} \geq \frac{7}{32}.$$

For $r = 2$, symmetry makes G descend in all slots to a two-dimensional quotient. Lemma IV.48.6 gives the same bound.

For sharpness, take $G = u^{\otimes 3} + v^{\otimes 3}$ with u, v independent. The two products equal one with probability $1/8$ each and simultaneously with probability $1/64$. Their sum is nonzero with probability $1/8 + 1/8 - 2/64 = 7/32$. $\square$

IV.48.3 A uniform higher-degree consequence

Theorem IV.48.8 (Non-pure support in arbitrary degree). *For $s \geq 3$, every symmetric non-pure map $F : V^s \rightarrow W$ satisfies*

$$\mathbb{P}[F \neq 0] \geq \frac{7}{2^{s+2}}. \quad (\text{IV.48.4})$$

Proof. Use Lemma IV.48.3 repeatedly to obtain a non-pure trilinear contraction after fixing $s - 3$ inputs. Its support is at least $7/32$ by Theorem IV.48.7. Lift support back $s - 3$ times using Lemma IV.48.4: $\mathbb{P}[F \neq 0] \geq 2^{-(s-3)}(7/32) = 7/2^{s+2}$. $\square$

Sharpness is established for $s = 3$, not claimed for every $s > 3$. The improvement over the elementary nonzero-map bound 2^{-s} is the factor $7/4$, and it uses symmetry and the exclusion of pure tensors.

IV.49 Universal thresholds and exact phase removal

Source 47, delivered Section 8.

IV.49.1 The noncanonical alternative

Theorem IV.49.1 (A quantitative obstruction dichotomy). *Let T be symmetric and nonintegrable and $\|f\|_\infty \leq 1$. Either $T = C_d^{u,v} + S$ for independent u, v and integrable S , in which case $\mathcal{M}(T) = b_d$, or*

$$\mathcal{E}_T(f) \leq \begin{cases} 1/4, & d = 3, \\ 1 - 6/2^d, & d = 4, 5, \\ 1 - 7/2^d, & d \geq 6. \end{cases} \quad (\text{IV.49.1})$$

Proof. For $d = 3$, if the nonzero alternating defect has rank two it is $u \otimes v + v \otimes u$, so subtracting $C_3^{u,v}$ leaves an integrable tensor. Otherwise its rank is at least four, and the shear profile gives $\mathcal{E}_T(f) \leq 1/4$.

Let $d \geq 4$ and $s = d - 3$. If Φ_T is pure, Theorem IV.47.4 gives the canonical alternative. Otherwise it is non-pure. For $s = 1$ it is a linear map of rank at least two, so its support has measure at least $3/4$. For $s = 2$ the support is at least $3/8$ by Proposition IV.48.5. For $s \geq 3$ it is at least $7/2^{s+2}$ by Theorem IV.48.8. Apply (IV.44.7) in the three cases to obtain, respectively, $5/8$, $13/16$, and $1 - 7/2^{s+3} = 1 - 7/2^d$. $\square$

Proof of Theorem F.52.2. The lower bound is the canonical example. In degree three the two alternatives in Theorem IV.49.1 have upper bounds $1/2$ and $1/4$; in degree four they have $11/16$ and $5/8$. In degree five both are at most $13/16$, and in degree six both are at most $57/64$. For

$d \geq 7$, the noncanonical upper bound $1 - 7/2^d$ is at least the canonical value $1 - (d + 1)/2^d$. This proves (F.52.6) and the four exact values.

Energy strictly above the universal nonintegrable upper bound forces $T \in \mathcal{I}_d(V)$. The explicit primitive exists by Theorem IV.43.3, and gauge invariance with $T = 0$ gives the lossless identity. The tensors C_d , $d = 3, 4, 5, 6$, with $f = 1$ are nonintegrable and attain the stated thresholds, proving sharpness. $\square$

IV.49.2 Where the sextic constant comes from

For a sextic tensor the frozen defect is trilinear. The two routes to $57/64$ are genuinely different. A pure defect is repaired to the canonical tensor and is handled by the derivative-period induction. A non-pure defect is nonzero on at least $7/32$ of the frozen triples, and each such triple costs at least $1/2$ in the rank profile. The total loss is therefore at least $7/64$.

This explains why merely repeating the earlier nonzero-map support estimate would not suffice. That estimate gives support only $1/8$ and hence loss $1/16$, leaving the old upper bound $15/16$. Even the simpler non-pure bilinear contraction argument yields loss only $6/64$. The sharp trilinear support census supplies the additional $1/64$ needed for the exact sextic result.

IV.49.3 Strict margins, not conclusions at equality

For $d = 4, 5, 6$, respectively, the implications are

$$\begin{aligned}\mathcal{E}_T(f) > 11/16 &\implies T \in \mathcal{I}_4(V), \\ \mathcal{E}_T(f) > 13/16 &\implies T \in \mathcal{I}_5(V), \\ \mathcal{E}_T(f) > 57/64 &\implies T \in \mathcal{I}_6(V).\end{aligned}$$

The assumption of symmetry includes the final frequency-evaluation slot. Without that assumption these are not established criteria. At equality there are nonintegrable examples, so replacing $>$ by $\geq$ is false. Neither the quintic nor the sextic proof classifies all endpoint tensors or all maximizing functions; the canonical examples prove sharpness without asserting uniqueness.

IV.50 Deficits, noisy selectors, and retained constraints

Source 47, delivered Section 9.

IV.50.1 A nonnegative canonical deficit recursion

For $d \geq 4$, let $m = d - 1$ and $g_z = \partial_z f$. On $u(z) = 1$ put

$$\tilde{g}_z = g_z e(-v(z)P_{u,m}).$$

Use conditional normalized averages over the indicated nonempty sets to define

$$\begin{aligned}D_{00} &= \mathbb{E}_{z:u(z)=v(z)=0}(1 - \mathcal{E}_0(g_z)), \\ D_{01} &= \mathbb{E}_{z:u(z)=0, v(z)=1}(1 - 2^{1-m} - \mathcal{E}_{U_m}(g_z)), \\ D_1 &= \mathbb{E}_{z:u(z)=1}(b_m - \mathcal{E}_{C_m}(\tilde{g}_z)).\end{aligned}\tag{IV.50.1}$$

Proposition IV.50.1 (Exact deficit decomposition). *Each quantity in (IV.50.1) is nonnegative, and*

$$b_d - \mathcal{E}_{C_d}(f) = \frac{1}{4}D_{00} + \frac{1}{4}D_{01} + \frac{1}{2}D_1. \quad (\text{IV.50.2})$$

In particular, if $\mathcal{E}_{C_d}(f) \geq b_d - \varepsilon$, then $D_{00}, D_{01} \leq 4\varepsilon$ and $D_1 \leq 2\varepsilon$.

Proof. Nonnegativity is exactly the three bounds in the proof of Theorem F.52.1. Subtract the exact sliced energy from the matching sum of slice caps in (IV.46.4). Gauge invariance on $u(z) = 1$ gives (IV.50.2). The last assertions follow because all terms are nonnegative. $\square$

On the $(0, 1)$ quarter, let $a_{z, \mathbf{h}}$ be the squared selected coefficient in the definition of $\mathcal{E}_{U_m}(g_z)$, and let A be the set where not all $u(h_i)$ equal one. The forbidden coefficients are exactly zero, so

$$D_{01} = \mathbb{E}_{z:u=0, v=1} \mathbb{E}_{\mathbf{h}} \mathbf{1}_A (1 - a_{z, \mathbf{h}}).$$

Thus for $0 < \tau \leq 1$, under $\mathcal{E}_{C_d}(f) \geq b_d - \varepsilon$,

$$\mathbb{P}[a_{z, \mathbf{h}} \leq 1 - \tau \mid u(z) = 0, v(z) = 1, \mathbf{h} \in A] \leq \frac{4\varepsilon}{(1 - 2^{1-m})\tau}. \quad (\text{IV.50.3})$$

This is a dimension-independent necessary condition on near-extremizers. It does not yet reconstruct a single phase that simultaneously accounts for all the nearly extremal slices.

IV.50.2 Corrupted frequency selectors

For an arbitrary selector $\rho : V^{d-1} \rightarrow V^*$, define $\mathcal{E}_\rho(f)$ by replacing $T(\mathbf{h}, \cdot)$ with $\rho(\mathbf{h})$ in (F.52.1). Suppose ρ agrees with the selector of a symmetric multilinear T outside a set of measure at most η .

Proposition IV.50.2 (Robust sharp certification). *For $\|f\|_\infty \leq 1$,*

$$|\mathcal{E}_\rho(f) - \mathcal{E}_T(f)| \leq \eta. \quad (\text{IV.50.4})$$

Put $\theta_d = 1 - \min\{d + 1, 7\}/2^d$. If a certified lower bound L for $\mathcal{E}_\rho(f)$ satisfies $L - \eta > \theta_d$, then T is integrable and

$$\|fe(-P_T)\|_{U_d}^{2^d} \geq L - \eta. \quad (\text{IV.50.5})$$

For $d = 3, 4, 5, 6$ the underlying uncorrupted thresholds are sharp.

Proof. Every squared coefficient lies in $[0, 1]$. The summands agree outside the exceptional set, proving (IV.50.4). The universal threshold and exact phase removal then give (IV.50.5). $\square$

A numerical error budget must be subtracted in forming the certified lower bound L . The proposition assumes an actual symmetric multilinear model; it does not construct one from noisy observations.

IV.50.3 Approximate periods with an explicit error term

Lemma IV.50.3 (Stable forbidden-frequency bound). *For a complex function G , a shift w , and a character $\chi_\xi(x) = (-1)^{\xi(x)}$ with $\xi(w) = 1$,*

$$|\widehat{G}(\xi)|^2 \leq \frac{1}{4} \|G - \tau_w G\|_2^2. \quad (\text{IV.50.6})$$

Proof. Translation gives $2\widehat{G}(\xi) = \mathbb{E}_x(G(x) - G(x+w))\chi_\xi(x)$. Apply Cauchy–Schwarz. $\square$

Consequently, in a family of one-bounded test functions with a set B of nominally forbidden coefficients, if a suitable period error is at most δ_b for each $b \in B$, then

$$\text{average squared coefficient} \leq 1 - \mathbb{P}(B) + \frac{1}{4} \mathbb{E}_b \mathbf{1}_B \delta_b^2. \quad (\text{IV.50.7})$$

The exact derivative theorem is the case $\delta_b = 0$. This formulation keeps analytic perturbations distinct from corruption of the selector, which is controlled by (IV.50.4) instead.

IV.51 Sources 65 and 69: Boolean phase integration in every degree [write]

Added in the eleventh write, 7 October 2026.

The next fifteen sections are source 65’s Sections 2–8 and then source 69’s Sections 2–9, on Gowers’s phase-removal step (Proposition 17.2) in characteristic two. *Logically they continue source 47’s chapter* (Section IV.42), whose degree-seven and all-degree questions they answer, and the binary chapter of sources 46 and 43 (Section IV.25); they are printed here because no number printed earlier may move. Their introductions are Sections F.70 and F.74.

The common theorem, proved twice. For every $d \geq 3$ and $n \geq 2$ the largest selected energy of a nonintegrable symmetric d -linear tensor on $\mathbb{F}_2^n$ is exactly $1 - (d+1)/2^d$, every extremal tensor is the canonical $C_d^{u,v}$ modulo an integrable tensor, and energy above $1 - (d+1)/2^d$ certifies integrability, with lossless phase removal (source 65’s Theorem F.70.1; source 69’s Theorem F.74.1, with the classification for $d \geq 4$ and the cubic law for $d = 3$). Every other nonintegrable tensor has energy at most $1 - 3d/2^{d+1}$ (both), and at most $1 - (1 - 2^{-r})d/2^{d-1}$ when the radical of integrable contractions has codimension r (source 65’s Theorem F.70.2). The key in both is that the directions with integrable contractions form exactly the radical of the repeated-variable defect (source 65’s K_T , source 69’s $R(T)$; note the different letters), so that a noncanonical tensor keeps at least three quarters of its contractions nonintegrable and induction on the degree closes. Source 47 had proved the canonical maximum in every degree and the universal value only through degree six ($1 - \min\{d+1, 7\}/2^d$); its Research questions V.176.1 and V.176.2 and the tensor part of V.176.4 are answered (item IV.11): the septic value is $15/16$. The two sources were written independently on the same day and neither cites the other; source 65 is printed first because its statement covers $d = 3$ and carries the radical-sensitive cap, and source 69 is a credited second route, printed whole with its further results. The integration criterion is Tao–Ziegler’s and Tidor’s, the cubic law and canonical family sources 43 and 47’s, credited by both and reproved.

Further results. Source 65: complete repair geometry ($T|_{H^d}$ integrable iff H lies in $\ker u$, $\ker v$ or $\ker(u+v)$, Theorem IV.57.2), the exact number of extremal classes, $(2^n - 1)(2^n - 2)/6$ for $d = 3$ and $(2^n - 1)(2^{n-1} - 1)$ for $d \geq 4$ (Theorem IV.57.4), dimension-independent amplitude stability

(Proposition IV.57.5), robust certification and a retained contraction profile (Theorem IV.58.1, Proposition IV.58.3). Source 69: the extremal classes are the flags $\ell \subset \Pi$ in V^* , also $(2^n - 1)(2^{n-1} - 1)$ in number (Theorem IV.63.1; the counts agree, checked at the write), the exact minimal quotient $m_{\text{eff}}(T) = \text{codim } K(T) \leq \rho + \sum_{j=1}^{\min(\rho, d-2)} \binom{\rho}{j}$, sharp for every $d \geq 4$ and $\rho \geq 1$ (Theorems F.74.2, IV.64.1, IV.65.1), a fixed-radical normal form (Theorem IV.64.4) and energy bounds on the contraction rank (Corollary IV.66.1).

Relation to the formal statements (checked at HEAD fb9602e55). `lemma_17_1` and `proposition_17_2` (Sections17_18.lean, lines 38 and 48) have the exact companions `lemma_17_1_holds` (Proofs17Phase.lean, line 374) and `proposition_17_2_holds` (Proofs17PhaseRemoval.lean, line 373); they are stated on $\mathbb{Z}/N\mathbb{Z}$ with N prime, so characteristic two lies outside them, as both sources say. Nothing in this chapter is formalized.

IV.52 Finite calculus and the integration criterion

Source 65, delivered Section 2.

IV.52.1 Cube expansion and phase changes

Lemma IV.52.1 (Selected cube identity and slicing). *For a multilinear tensor T and an arbitrary complex function f ,*

$$\mathcal{E}_T(f) = \mathbb{E}_{x, h_1, \dots, h_d} \partial_{h_1} \cdots \partial_{h_d} f(x) (-1)^{T(h_1, \dots, h_d)}. \quad (\text{IV.52.1})$$

If T is symmetric and $d \geq 3$, then, with $T_z(h_1, \dots, h_{d-1}) = T(h_1, \dots, h_{d-1}, z)$,

$$\mathcal{E}_T(f) = \mathbb{E}_z \mathcal{E}_{T_z}(\partial_z f). \quad (\text{IV.52.2})$$

Proof. For a fixed tuple $\mathbf{h} = (h_1, \dots, h_{d-1})$ let $g = \partial_{h_1} \cdots \partial_{h_{d-1}} f$ and $\xi(x) = T(\mathbf{h}, x)$. Expanding the squared Fourier coefficient and changing one base point to the other plus y gives

$$\left| \mathbb{E}_x g(x) (-1)^{\xi(x)} \right|^2 = \mathbb{E}_{x, y} g(x+y) \overline{g(x)} (-1)^{\xi(y)}.$$

Average over $\mathbf{h}$ to obtain (IV.52.1). Multiplicative differences commute, as follows either from their four-factor expansion or from the cube product. Freeze the last direction as z , commute its difference onto f , and use symmetry to put the remaining evaluation slot in the required position. This gives (IV.52.2). $\square$

All these are identities of products and finite sums. They remain valid when f has zeros; no quotient by f is involved.

Lemma IV.52.2 (Gauge invariance). *If $S \in \mathcal{I}_d(V)$ and P is a primitive of S , then*

$$\mathcal{E}_{T+S}(fe(P)) = \mathcal{E}_T(f), \quad \mathcal{M}(T+S) = \mathcal{M}(T). \quad (\text{IV.52.3})$$

Proof. The top multiplicative derivative of $e(P)$ is $(-1)^S$. The two extra factors of $(-1)^S$ in (IV.52.1) cancel. Multiplication by $e(P)$ is a bijection of the unit polydisc of functions. $\square$

In particular an integrable tensor has maximum one, and gauge invariance with $T = 0$ gives the identity (F.70.6).

IV.52.2 A self-contained Boolean integration proof

Theorem IV.52.3 (Repeated-variable integration criterion). *For a symmetric d -linear tensor with $d \geq 3$,*

$$T \in \mathcal{I}_d(V) \iff \Phi_T = 0. \quad (\text{IV.52.4})$$

Every symmetric bilinear tensor is integrable. In a basis with coordinate bits $x_1, \dots, x_n$, an integrable tensor has an explicit primitive

$$P_T(x) = \sum_{\substack{\emptyset \neq A \subseteq [n] \\ |A| \leq d}} \frac{c_A \prod_{i \in A} |x_i|}{2^{d-|A|+1}} \pmod{1}, \quad (\text{IV.52.5})$$

where $|x_i| \in \{0, 1\}$ and c_A is the value of T on any length- d basis tuple whose set of occurring indices is A .

Proof. Translation by a Boolean direction a squares to the identity. Thus, on circle-valued functions,

$$\Delta_a^2 = -2\Delta_a, \quad \Delta_a^2 \Delta_b = \Delta_a \Delta_b^2. \quad (\text{IV.52.6})$$

Commutation and (IV.52.6) prove the necessary repeated-variable relation for the top derivative of any primitive.

Conversely suppose $\Phi_T = 0$. By symmetry a basis coefficient depends on its multiplicities. If both i and j occur and i occurs at least twice, the identity $T(e_i, e_i, e_j, \mathbf{z}) = T(e_i, e_j, e_j, \mathbf{z})$ transfers one occurrence from i to j . All positive multiplicity lists of total d on a given support are connected by these transfers. Hence c_A is well defined.

For $a = |A|$, consider the summand $P_A = 2^{-(d-a+1)} \prod_{i \in A} |x_i|$ in (IV.52.5). We use

$$\Delta_{e_i}^t = (-2)^{t-1} \Delta_{e_i} \quad (t \geq 1), \quad \Delta_h = \prod_{i: h_i=1} (I + \Delta_{e_i}) - I.$$

A term in the expansion of $d+1$ differences has at least $d+1$ coordinate differences. If a coordinate outside A occurs, the term vanishes. Otherwise, with $r \leq a$ distinct coordinates among $t \geq d+1$ differences, its numerator is divisible by 2^{t-r} and therefore by 2^{d-a+1} . It is zero modulo one. This proves degree at most d .

For exactly d coordinate differences, a support properly contained in A again gives an integer, and a coordinate outside A gives zero. On a tuple with support exactly A , the value is $\pm 1/2 = 1/2$ modulo one, independent of the base point. Thus (IV.52.5) has the required basis coefficients in its top derivative. A top derivative of a degree-at-most- d polynomial is independent of the base point. It is also multiadditive, because the error term in $\Delta_{a+b} = \Delta_a + \Delta_b + \Delta_a \Delta_b$ is an extra difference. Agreement on basis tuples proves (F.70.2). For $d = 2$ no multiplicity transfer is needed and the same construction applies to every symmetric tensor. $\square$

Corollary IV.52.4 (Dimension of the integrable space). *The dimension over $\mathbb{F}_2$ is*

$$q_d(n) := \dim \mathcal{I}_d(V) = \sum_{j=1}^{\min(d,n)} \binom{n}{j}. \quad (\text{IV.52.7})$$

For a nonzero $u \in V^*$ the pure tensor $U_m = u^{\otimes m}$ has primitive $P_{u,m}(x) = |u(x)|/2^m$.

Proof. The independent bits c_A parameterize all integrable tensors uniquely. For the final assertion apply the one-coordinate construction and pull back along u . $\square$

IV.53 The cubic spectral base case

Source 65, delivered Section 3.

The only spectral input needed for the new all-degree induction is the cubic estimate. This section supplies a complete finite-dimensional proof of the estimate used in the earlier drafts.

IV.53.1 Projective translations

Let $B : V \times V \rightarrow \mathbb{F}_2$ be bilinear and set $\Omega_B(a, b) = B(a, b) + B(b, a)$. This is alternating. Its rank is even: if it is nonzero, choose a, b with $\Omega_B(a, b) = 1$, split off the nondegenerate two-dimensional span of a, b , and repeat on its orthogonal complement. Write $\text{rank } \Omega_B = 2s$.

For complex g, h define

$$Q_B(g, h) = \mathbb{E}_a \left| \mathbb{E}_x g(x+a) \overline{h(x)} (-1)^{B(a,x)} \right|^2, \quad Q_B(g) = Q_B(g, g). \quad (\text{IV.53.1})$$

Lemma IV.53.1 (Finite projective orthogonality). *For arbitrary complex g, h ,*

$$Q_B(g, h) \leq 2^{-s} \|g\|_2^2 \|h\|_2^2, \quad Q_B(g) = \mathbb{E}_{x,a,b} \partial_a \partial_b g(x) (-1)^{B(a,b)}. \quad (\text{IV.53.2})$$

Proof. Use the normalized inner product $\langle g, h \rangle = \mathbb{E}_x g(x) \overline{h(x)}$, linear in its first argument. The operators

$$(U_a g)(x) = (-1)^{B(a,x)} g(x+a)$$

are unitary and satisfy

$$U_a U_b = (-1)^{B(b,a)} U_{a+b}, \quad U_a U_b = (-1)^{\Omega_B(a,b)} U_b U_a.$$

The positive operator

$$M_g v = \mathbb{E}_a \langle v, U_a g \rangle U_a g$$

has trace $\|g\|_2^2$ and commutes with every U_b : conjugation permutes its rank-one summands, and the scalar signs cancel.

Choose a maximal isotropic subspace H containing the radical of Ω_B . The elementary symplectic decomposition above gives $H = H^\perp$ and $[V : H] = 2^s$. Every nonzero subspace invariant under all U_a has dimension at least 2^s . To see this, simultaneously diagonalize the commuting unitaries U_h , $h \in H$, on that invariant subspace and choose a nonzero joint eigenvector w . The vectors $U_a w$ for representatives of V/H have distinct joint eigenvalues: their ratios on H are $h \mapsto (-1)^{\Omega_B(h,a)}$, and distinct cosets give distinct characters since $H^\perp = H$. Joint eigenvectors of commuting unitaries with distinct joint eigenvalues are orthogonal.

Each positive eigenspace of M_g is invariant under all U_a , hence has dimension at least 2^s . Positivity and the trace identity imply $\|M_g\|_{\text{op}} \leq 2^{-s} \|g\|_2^2$. Since $Q_B(g, h) = \langle M_g h, h \rangle$, this proves the inequality.

For the second identity expand the square in (IV.53.1), and replace one base point by the other plus b . The character factor becomes $(-1)^{B(a,b)}$. The product of the four function values is the conjugate of $\partial_a \partial_b g(x)$. The resulting cube average is real, because it equals the real number $Q_B(g)$ after conjugation, so the stated identity follows. $\square$

IV.53.2 The exact shear

Theorem IV.53.2 (Exact cubic obstruction law). *For every symmetric trilinear T ,*

$$\mathcal{M}(T) = 2^{-\text{rank } \Omega_T/2}. \quad (\text{IV.53.3})$$

In particular a nonintegrable cubic has maximum at most $1/2$, and it has maximum $1/2$ exactly when it is canonical modulo integration.

Proof. The map $a \mapsto T(a, a, b)$ is additive: the two cross terms in the expansion at $a + a'$ cancel by symmetry. It follows that Ω_T is alternating bilinear. The following identity is valid pointwise for every complex f :

$$\partial_a \partial_{a+w} f = \overline{\partial_a (f \tau_w f)}, \quad \tau_w f(x) = f(x + w). \quad (\text{IV.53.4})$$

Indeed both sides are the product $f(x)f(x+w)\overline{f(x+a)f(x+a+w)}$. In the cubic average (IV.52.1), substitute $c = a + w$. The tensor factor becomes the bilinear form

$$B_w(a, b) = T(a, a, b) + T(a, b, w).$$

Its alternating part is Ω_T , since the term containing w is symmetric in a, b . Apply the remaining difference to (IV.53.4) and use the real self-cube identity in Lemma IV.53.1. This gives

$$\mathcal{E}_T(f) = \mathbb{E}_w Q_{B_w}(f \tau_w f) \leq 2^{-\text{rank } \Omega_T/2} \mathbb{E}_w \|f \tau_w f\|_2^4 \leq 2^{-\text{rank } \Omega_T/2}.$$

For equality, symplectic coordinates express a rank- $2s$ form as $\Omega_T = \sum_{j=1}^s (u_j \otimes v_j + v_j \otimes u_j)$, with the $2s$ coordinate forms independent. Direct calculation gives $\Omega_{C_3^{u_j, v_j}} = u_j \otimes v_j + v_j \otimes u_j$. Therefore $S = T - \sum_j C_3^{u_j, v_j}$ is integrable. At the constant function, the energy of one block $C_3^{u, v}$ is $1/2$: its selected frequency vanishes on exactly half of the pairs of input vectors. Independent blocks give independent selected-frequency conditions, so the energy of their sum at $f = 1$ is 2^{-s} . Gauge invariance supplies an extremizer for T . For $s = 0$ the tensor is integrable and the value is one. For $s = 1$ this also proves precisely the final canonical characterization. $\square$

The block calculation concerns an explicit extremizer. It is not an argument that arbitrary optimizing functions factor as products; the spectral upper bound already handled every bounded function.

IV.54 The canonical energy in every degree

Source 65, delivered Section 4.

The all-degree canonical result is an ingredient from [264], not a new claim of this manuscript. Its short proof is included because the universal theorem needs its exact value, including stability under adding coordinates that the tensor does not use.

IV.54.1 Exact derivative periods

Lemma IV.54.1 (Even-direction periods). *For $r \geq 1$, $F = \partial_{h_1} \cdots \partial_{h_r} f$ satisfies*

$$F(x + h_i) = \overline{F(x)} \quad (1 \leq i \leq r), \quad F(x + w) = F(x) \quad \left(w \in \left\{ \sum_i \varepsilon_i h_i : \sum_i \varepsilon_i = 0 \right\} \right). \quad (\text{IV.54.1})$$

Consequently $\hat{F}(\xi) \neq 0$ implies $\xi(h_1) = \cdots = \xi(h_r)$.

Proof. The cube product for $F(x)$ has one labelled factor at each Boolean vertex, with conjugation parity $r - |\varepsilon|$. Translation by h_i toggles the corresponding vertex coordinate, reversing conjugation parity. This conjugates the full product. Two such translations preserve it; their sums generate the even-direction space in (IV.54.1). If a character is nontrivial on one of these periods, translation changes the sign of its Fourier coefficient without changing the transformed function. The coefficient is therefore zero. This proof also covers coincident vertices and zeros of f . $\square$

Lemma IV.54.2 (A pure tensor on a constrained derivative). *Let $u \neq 0$, $m \geq 2$, and $u(z) = 0$. Then*

$$\mathcal{E}_{U_m}(\partial_z f) \leq 1 - 2^{1-m}, \quad U_m = u^{\otimes m}, \quad \|f\|_\infty \leq 1. \quad (\text{IV.54.2})$$

The constant function $f = 1$ attains equality.

Proof. In the selected-energy test for U_m , the selected frequency is u exactly when all $m - 1$ sampled directions satisfy $u(h_i) = 1$; otherwise it is zero. On that first set the transformed function is $\partial_{h_1} \cdots \partial_{h_{m-1}} \partial_z f$. It has period $h_1 + z$, but $u(h_1 + z) = 1$. The selected coefficient vanishes by Lemma IV.54.1. The excluded set has measure $2^{-(m-1)}$; all remaining squared coefficients are at most one. For $f = 1$ they equal one exactly when the selected frequency is zero, giving equality. $\square$

Although U_m integrates and has unrestricted maximum one, the fixed slice $\partial_z f$ cannot attain that maximum when $u(z) = 0$. This derivative constraint is the essential extra loss in the canonical induction.

IV.54.2 The sharp canonical formula

Proposition IV.54.3 (Canonical defect and constant energy). *For $d \geq 3$ and independent u, v ,*

$$\Phi_{C_d^{u,v}}(z)(a, b) = \left(\prod_{j=1}^{d-3} u(z_j) \right) (u(a)v(b) + v(a)u(b)), \quad \mathcal{E}_{C_d^{u,v}}(1) = b_d. \quad (\text{IV.54.3})$$

In particular $C_d^{u,v}$ is nonintegrable.

Proof. In the two repeated-variable evaluations the terms placing v in either repeated slot cancel, as do the terms placing v in a frozen slot. The two remaining terms give the displayed defect. It is nonzero.

Set $k = d - 1$. The selected frequency for the constant function has the form $Av + Bu$, with

$$A = \prod_{j=1}^k u(h_j), \quad B = \sum_{i=1}^k v(h_i) \prod_{j \neq i} u(h_j).$$

It is zero when at least two of the $u(h_j)$ vanish. When exactly one vanishes, $A = 0$ and B is its independent unbiased v bit. When none vanishes, $A = 1$. The probability of a nonzero selected frequency is $2^{-k} + k2^{-k}/2 = (d + 1)/2^d$, proving the energy formula. $\square$

Theorem IV.54.4 (Canonical maximum, with arbitrary inactive coordinates). *For every $d \geq 3$, independent $u, v \in V^*$, and $S \in \mathcal{I}_d(V)$,*

$$\mathcal{M}(C_d^{u,v} + S) = b_d. \quad (\text{IV.54.4})$$

There is no assumption that the tested function factors through (u, v) .

Proof. Gauge invariance reduces to $S = 0$. The cubic case is Theorem IV.53.2. Suppose the result holds in degree $m = d - 1$. Directly from the definition,

$$(C_d^{u,v})_z = u(z)C_m^{u,v} + v(z)U_m.$$

Slice the energy as in (IV.52.2). The four pairs $(u(z), v(z))$ occur with probability $1/4$ each. On $(0, 0)$ the slice cap is one. On $(0, 1)$ use the derivative-constrained cap $1 - 2^{1-m}$ from Lemma IV.54.2. On the half where $u(z) = 1$, the slice tensor is C_m plus an integrable tensor, so gauge invariance and the induction hypothesis give cap b_m . Consequently

$$\begin{aligned} \mathcal{E}_{C_d}(f) &\leq \frac{1}{4} + \frac{1}{4}(1 - 2^{2-d}) + \frac{1}{2}b_{d-1} \\ &= 1 - \frac{d+1}{2^d} = b_d. \end{aligned} \tag{IV.54.5}$$

Proposition IV.54.3 supplies equality. Every sliced function, including a phase-twisted one, was allowed to depend on all of V , so no quotient-factorization assumption has entered the argument. $\square$

IV.55 Defect geometry and the integrable-slice identity

Source 65, delivered Section 5.

This section isolates the algebra that turns the canonical theorem into a universal theorem. The distinction between a defect's image and its first-slot radical is indispensable.

IV.55.1 Alternation and compatibility

Lemma IV.55.1 (Defect compatibility). *For $d \geq 3$, Φ_T is a symmetric $(d-3)$ -linear map to $\text{Alt}^2(V)$. For $d \geq 4$ and a tuple $\mathbf{w} \in V^{d-4}$,*

$$\Phi_T(c, \mathbf{w})(a, b) + \Phi_T(b, \mathbf{w})(a, c) + \Phi_T(a, \mathbf{w})(b, c) = 0. \tag{IV.55.1}$$

Proof. Additivity in a repeated variable follows because its two cross terms cancel in characteristic two. The value at $a = b$ is zero. The other multilinearity and symmetry assertions follow immediately from those of T . Expanding (IV.55.1) yields six terms of T , which pair equally by symmetry and cancel. $\square$

Lemma IV.55.2 (A one-form compatibility identity). *If $u \neq 0$ and $\Omega \in \text{Alt}^2(V)$ satisfy*

$$u(c)\Omega(a, b) + u(b)\Omega(a, c) + u(a)\Omega(b, c) = 0 \quad \text{for all } a, b, c,$$

then $\Omega = u \otimes v + v \otimes u$ for some $v \in V^$. When $\Omega \neq 0$, u, v are independent and $\text{rank } \Omega = 2$.*

Proof. Choose c_0 with $u(c_0) = 1$ and put $v(x) = \Omega(x, c_0)$. The identity with $c = c_0$ gives the formula. Its nonzero case has rank two by an immediate two-coordinate calculation. $\square$

IV.55.2 Radical codimension one is the canonical case

Proposition IV.55.3 (Canonical recognition). *For symmetric nonintegrable T of degree $d \geq 4$, the following are equivalent:*

- (i) $r(T) = 1$;
- (ii) $\Phi_T(\mathbf{z}) = \Omega \prod_j u(z_j)$ for nonzero u and nonzero $\Omega \in \text{Alt}^2(V)$;
- (iii) $T = C_d^{u,v} + S$ for independent u, v and $S \in \mathcal{I}_d(V)$.

Proof. By symmetry of Φ_T , its first-slot radical is a radical in every frozen slot. If it has codimension one, the map descends in each frozen slot to V/K_T , which is one-dimensional. Thus it has the form in (ii). Conversely such a nonzero map has first-slot radical $\ker u$.

Set every frozen variable other than the one displayed in (IV.55.1) equal to a vector on which u is one. The identity becomes the hypothesis of Lemma IV.55.2. Hence $\Omega = u \otimes v + v \otimes u$ with u, v independent. By (IV.54.3), $T - C_d^{u,v}$ has zero defect and is integrable. This proves (ii) $\Rightarrow$ (iii). The reverse implication follows from that same defect formula and the integration criterion. $\square$

This recognition result is the pure-defect classification of [264], expressed in the radical notation needed for the next lemma. A map with one-dimensional *image* need not have radical codimension one; no inference of that kind is used here.

IV.55.3 The exact set of integrable contractions

Lemma IV.55.4 (Integrable-slice identity). *For $d \geq 4$ and $T \in \text{SymMult}^d(V)$,*

$$\{z \in V : T_z \in \mathcal{I}_{d-1}(V)\} = K_T, \quad \mathbb{P}_z[T_z \in \mathcal{I}_{d-1}(V)] = 2^{-r(T)}. \quad (\text{IV.55.2})$$

Here $K_T = V$ and $r(T) = 0$ when T is integrable.

Proof. For a fixed direction z ,

$$\begin{aligned} \Phi_{T_z}(z_2, \dots, z_{d-3})(a, b) &= T(a, a, b, z_2, \dots, z_{d-3}, z) \\ &\quad + T(a, b, b, z_2, \dots, z_{d-3}, z) \\ &= \Phi_T(z, z_2, \dots, z_{d-3})(a, b). \end{aligned}$$

The equivalence in (IV.55.2) is exactly Theorem IV.52.3 in degree $d - 1$. This includes $d = 4$, when the slice is cubic and its defect has no frozen variables. The map $z \mapsto \Phi_T(z, \cdot)$ is linear, so K_T is a subspace of size $2^{n-r(T)}$. Dividing by $|V|$ proves the probability formula. $\square$

Proposition IV.55.5 (One-step transfer of a universal bound). *Let $d \geq 4$. Suppose every nonintegrable symmetric $(d - 1)$ -linear tensor has maximum at most B . Then every symmetric d -linear tensor satisfies*

$$\mathcal{E}_T(f) \leq 2^{-r(T)} + (1 - 2^{-r(T)})B, \quad \|f\|_\infty \leq 1. \quad (\text{IV.55.3})$$

Proof. Use exact slicing (IV.52.2). When $z \in K_T$, the slice energy is at most one. When $z \notin K_T$, the slice tensor is nonintegrable by Lemma IV.55.4, and $\|\partial_z f\|_\infty \leq 1$, so its energy is at most B . Average and use the exact slice proportions. $\square$

No replacement of the common function by independent functions is made in the identity. The inequality subsequently relaxes each nonintegrable slice to an arbitrary bounded function, which is legitimate for an upper bound. The canonical case, where this relaxation would lose too much, was already treated separately in Section IV.54.

IV.56 Closing the induction in every degree

Source 65, delivered Section 6.

Proof of Theorems F.70.1 and F.70.2. The cubic theorem supplies the base $b_3 = 1/2$ and the cubic endpoint classification. Fix $d \geq 4$ and assume that every nonintegrable symmetric tensor of degree $d-1$ has maximum at most b_{d-1} .

Let T be a nonintegrable symmetric d -linear tensor. Its defect is nonzero, so $r(T) \geq 1$. If $r(T) = 1$, canonical recognition and Theorem IV.54.4 give $\mathcal{M}(T) = b_d$.

If $r = r(T) \geq 2$, Proposition IV.55.5 with $B = b_{d-1}$ gives

$$\begin{aligned} \mathcal{M}(T) &\leq 2^{-r} + (1 - 2^{-r}) \left(1 - \frac{d}{2^{d-1}}\right) \\ &= 1 - (1 - 2^{-r}) \frac{d}{2^{d-1}} \\ &\leq 1 - \frac{3d}{2^{d+1}} = b_d - \frac{d-2}{2^{d+1}} < b_d. \end{aligned} \tag{IV.56.1}$$

This proves the universal upper bound and the stated radical-sensitive cap. The canonical example gives the matching lower bound. Induction therefore proves the universal formula for all degrees.

The strict final inequality in (IV.56.1) also proves the endpoint classification: a noncanonical tensor cannot attain b_d . Conversely every canonical tensor modulo integration does attain b_d . The cubic case was already classified by its alternating rank. Finally, energy strictly above b_d excludes nonintegrability; the primitive and exact phase-removal identity follow from the integration criterion and gauge invariance. At $T = C_d$, $f = 1$ the energy is exactly b_d and T does not integrate, so the threshold is strict and sharp. $\square$

IV.56.1 Why the prior support bottleneck disappears

A direct estimate based on nonzero values of the repeated defect can forget how many *whole contractions* remain nonintegrable. The preceding manuscript propagated a fixed trilinear support coefficient through higher frozen degree. That yielded the loss $7/2^d$ beyond degree six. The transfer in (IV.55.3) instead propagates the *sharp degree-dependent extremal loss* $d/2^{d-1}$. The distinction is decisive: the factor $3/4$ beats the target loss, since

$$\frac{3}{4} \frac{d}{2^{d-1}} - \frac{d+1}{2^d} = \frac{d-2}{2^{d+1}} > 0.$$

There is consequently no need to solve the separate problem of optimal non-pure multilinear support in each degree. That problem remains interesting but is not a missing hypothesis of the theorem.

IV.56.2 Septic and endpoint consequences explicitly

For $d = 7$ the prior interval $[120/128, 121/128]$ collapses to

$$M_{7,n}^{\text{sym}} = 120/128 = 15/16 \quad (n \geq 2).$$

A noncanonical septic obeys the stronger cap $235/256$, whereas the extremal value is $240/256$. For quintics the old noncanonical cap $13/16$ is replaced by $49/64 < 52/64$; for sextics the old cap $57/64$ is replaced by $55/64$. Thus the previously undecided noncanonical endpoint branches are excluded without an equality analysis of the old support estimates.

IV.57 Rigidity, repair geometry, and exact counting

Source 65, delivered Section 7.

IV.57.1 A finite rigidity window

Set

$$\gamma_3 = \frac{1}{4}, \quad \gamma_d = \frac{d-2}{2^{d+1}} \quad (d \geq 4), \quad \eta_d = b_d - \gamma_d. \quad (\text{IV.57.1})$$

Corollary IV.57.1 (Exact tensor rigidity near the endpoint). *If T is symmetric and $\mathcal{E}_T(f) > \eta_d$, then either T is integrable or T is canonical modulo integration. In particular, if T is known to be nonintegrable and*

$$\mathcal{E}_T(f) \geq b_d - \varepsilon, \quad 0 \leq \varepsilon < \gamma_d,$$

then $T = C_d^{u,v} + S$ exactly, with S integrable.

Proof. The noncanonical alternatives of Theorem F.70.2 have energy at most η_d . Exclude them. $\square$

This is exact rigidity of a discrete tensor class, not a claim that the function f is close to one particular phase. The strict window is a proved sufficient window; its optimal width is not established.

IV.57.2 All integrable restrictions of an endpoint tensor

Theorem IV.57.2 (Complete repair geometry). *Let $d \geq 3$ and $T = C_d^{u,v} + S$ with u, v independent and S integrable. For any subspace $H \leq V$,*

$$T|_{H^d} \text{ is integrable} \iff \dim \text{span}\{u|_H, v|_H\} \leq 1. \quad (\text{IV.57.2})$$

Equivalently, H is contained in at least one of

$$\ker u, \quad \ker v, \quad \ker(u + v). \quad (\text{IV.57.3})$$

These three distinct hyperplanes are exactly the maximal integrable restriction subspaces. The minimum repair codimension is one.

Proof. An integrable tensor remains integrable on restriction, so remove S . The restricted defect is

$$\left(\prod_j u|_H(z_j) \right) (u|_H(a)v|_H(b) + v|_H(a)u|_H(b)).$$

It vanishes when the two restricted forms are dependent. If they are independent, their alternating product is nonzero; for $d > 3$ choose each frozen direction with $u|_H(z_j) = 1$. The defect is then nonzero. For $d = 3$ the product is empty and the same conclusion follows directly. The integration criterion proves (IV.57.2).

Two forms over $\mathbb{F}_2$ are dependent precisely when $u|_H = 0$, $v|_H = 0$, or $(u + v)|_H = 0$. This proves the containment characterization. The three ambient kernels are distinct hyperplanes, all integrable on restriction. Every integrable subspace is contained in one of them, so these and only these are maximal. Codimension zero is impossible because T is nonintegrable; codimension one is achieved. $\square$

For $d \geq 4$, $\ker u$ is distinguished by the frozen radical of the defect. That does *not* make it the unique optimal repair hyperplane: there are always the three in (IV.57.3).

IV.57.3 Uniqueness of the recovered parameters

Proposition IV.57.3 (Canonical obstruction parameters). *For $d \geq 4$, the integrable equivalence class of $C_d^{u,v}$ uniquely determines u and the coset $v + \text{span}\{u\}$. Conversely these data uniquely determine its integrable equivalence class.*

Proof. The defect determines $K_T = \ker u$. Over $\mathbb{F}_2$ a hyperplane determines its nonzero defining form uniquely. Choose z_0 with $u(z_0) = 1$. Evaluation of all frozen slots at z_0 recovers $\Omega = u \otimes v + v \otimes u$. If $u \otimes w + w \otimes u = 0$, substituting a vector with $u = 1$ shows $w \in \text{span}\{u\}$. Thus Ω determines precisely the claimed coset of v .

Conversely, replacing v by $v + u$ changes $C_d^{u,v}$ by $(d \bmod 2)U_d$, which is integrable. Equal defects mean an integrable difference by Theorem IV.52.3. This proves both claims. $\square$

For cubics the distinguished u disappears: the class is determined by $u \wedge v$, hence by the two-dimensional subspace spanned by u, v . For higher degrees it is a two-dimensional subspace together with a specified nonzero one-dimensional subspace inside it.

Theorem IV.57.4 (Exact number of extremal tensor classes). *For $n \geq 2$, the number of nonintegrable classes modulo $\mathcal{I}_d(V)$ with maximum b_d is*

$$e_d(n) = \begin{cases} \frac{(2^n - 1)(2^n - 2)}{6}, & d = 3, \\ (2^n - 1)(2^{n-1} - 1), & d \geq 4. \end{cases} \quad (\text{IV.57.4})$$

The number of actual extremal tensors is $e_d(n)2^{q_d(n)}$, where $q_d(n)$ is given by (IV.52.7).

Proof. For $d \geq 4$ there are $2^n - 1$ choices for u . The quotient $V^*/\text{span}\{u\}$ has $2^{n-1} - 1$ nonzero cosets available for v . Proposition IV.57.3 makes these parameters injective and exhaustive.

For $d = 3$, a rank-two alternating form is $u \wedge v$ for an independent ordered pair. Its associated two-dimensional subspace has six ordered bases, and every invertible two-by-two matrix over $\mathbb{F}_2$ has determinant one, so all six bases give the same nonzero alternating form. Distinct such subspaces give distinct forms. This gives $(2^n - 1)(2^n - 2)/6$. Each integrable class contains exactly $2^{q_d(n)}$ tensors. $\square$

[write] **7 October 2026; sources 60–69.** *Against source 69* (checked at the write): $(2^n - 1)(2^{n-1} - 1)$ is the number of flags $\ell \subset \Pi$ of a line in a plane in V^* , source 69's count (its Theorem IV.63.1); at $d = 3$, $(2^n - 1)(2^n - 2)/6$ is the number of planes, the line not being determined by the class.

In dimension two there is one cubic extremal class and three extremal classes in each degree at least four. Since $q_d(2) = 3$, this gives eight cubic extremal tensors and twenty-four in every higher degree. In dimension three there are seven cubic classes and twenty-one higher-degree classes; quartics therefore have $21 \cdot 128 = 2688$ extremal tensors. These count *tensors admitting an extremizer*, not tensors whose extremizer must be the constant function.

IV.57.4 A necessary amplitude condition

Proposition IV.57.5 (Dimension-independent amplitude stability). *Let $d \geq 4$ and T be a canonical tensor modulo integration. If $\mathcal{E}_T(f) \geq b_d - \varepsilon$ and $\|f\|_\infty \leq 1$, then*

$$\mathbb{E}_x |f(x)|^2 \geq 1 - 4\varepsilon. \quad (\text{IV.57.5})$$

In particular every endpoint function for a nonintegrable tensor in these degrees has modulus one at every point.

Proof. Remove the integrable correction by a phase twist, preserving $|f|$. In the canonical induction the three slice caps are attained in sum exactly at b_d , and each slice lies below its cap. Therefore the contribution of the $(u(z), v(z)) = (0, 0)$ quarter gives

$$b_d - \mathcal{E}_{C_d}(f) \geq \frac{1}{4} \mathbb{E}_{z:u(z)=v(z)=0} (1 - \mathcal{E}_0(\partial_z f)).$$

For any one-bounded g and $m \geq 2$, the cube product, the triangle inequality, and arithmetic–geometric mean give

$$\mathcal{E}_0(g) \leq \mathbb{E}_x |g(x)|^{2^m} \leq \mathbb{E}_x |g(x)|^2.$$

Indeed apply arithmetic–geometric mean to the 2^m nonnegative vertex factors raised to the 2^m th power, then average; each vertex is uniform. For $g = \partial_z f$ the last mean is at most $\mathbb{E}_x |f(x)|^2$. Substitution proves (IV.57.5). At $\varepsilon = 0$, an average of numbers in $[0, 1]$ is one, so every modulus is one. $\square$

The coefficient four is not claimed optimal. This statement is about amplitudes only; phases can still have nontrivial freedom.

IV.58 Robust certification and retained contraction profiles

Source 65, delivered Section 8.

IV.58.1 Corrupted selectors and measurement error

For a selector $\rho : V^{d-1} \rightarrow V^*$, define $\mathcal{E}_\rho(f)$ by replacing $T(\mathbf{h}, \cdot)$ in (F.70.1) by $\rho(\mathbf{h})$.

Theorem IV.58.1 (Robust all-degree phase certification). *Suppose $\rho(\mathbf{h}) = T(\mathbf{h}, \cdot)$ outside a fraction δ of tuples, where T is symmetric and $d \geq 3$. Then*

$$|\mathcal{E}_\rho(f) - \mathcal{E}_T(f)| \leq \delta. \quad (\text{IV.58.1})$$

If L is a certified lower bound on $\mathcal{E}_\rho(f)$, then

$$L - \delta > b_d \implies T \in \mathcal{I}_d(V), \quad \|fe(-P_T)\|_{V^d}^{2^d} \geq L - \delta; \quad (\text{IV.58.2})$$

$$L - \delta > \eta_d \implies T \text{ is integrable or canonical modulo integration.} \quad (\text{IV.58.3})$$

Proof. Each squared coefficient lies in $[0, 1]$. The two summands agree on the uncorrupted tuples, proving (IV.58.1). Apply the sharp threshold and the rigidity window to $\mathcal{E}_T(f) \geq L - \delta$. $\square$

The theorem assumes an actual symmetric tensor model. It does not construct or verify that model from noisy data. Any numerical error in an energy estimate must already have been subtracted when forming L .

Lemma IV.58.2 (Perturbing the function). *For one-bounded f, g and any multilinear T ,*

$$|\mathcal{E}_T(f) - \mathcal{E}_T(g)| \leq 2^d \|f - g\|_1. \quad (\text{IV.58.4})$$

Proof. In (IV.52.1), telescope the difference of the two products over its 2^d labelled vertices. Every other factor has modulus at most one; the factor being replaced contributes $|f - g|$ at that vertex. Each vertex is uniform after averaging over x , even when some vertices coincide. Summing gives the bound. $\square$

For a test performed on g rather than f , subtract the additional $2^d \|f - g\|_1$ before applying the certification thresholds. No optimality of this perturbation constant is asserted.

IV.58.2 A stronger computable profile

Define a tensor cap $A_m(S)$ for symmetric tensors of degree $m \geq 3$ by

$$A_m(S) = \begin{cases} 1, & S \in \mathcal{I}_m(V), \\ 2^{-\text{rank } \Omega_S/2}, & m = 3, \\ b_m, & m \geq 4, r(S) = 1, \\ 1 - (1 - 2^{-r(S)})m/2^{m-1}, & m \geq 4, r(S) \geq 2. \end{cases} \quad (\text{IV.58.5})$$

The cases in the displayed order remove any overlap at an integrable cubic. The main theorems prove $\mathcal{M}(S) \leq A_m(S)$.

Proposition IV.58.3 (Retained contraction profile). *For $d \geq 4$,*

$$\mathcal{E}_T(f) \leq \mathbb{E}_z A_{d-1}(T_z). \quad (\text{IV.58.6})$$

For nonintegrable T of degree $d \geq 5$, let

$$\mu_T = \mathbb{P}_z[T_z \text{ is nonintegrable and noncanonical}].$$

Then

$$\mathcal{M}(T) \leq 1 - (1 - 2^{-r(T)}) \frac{d}{2^{d-1}} - \mu_T \gamma_{d-1}. \quad (\text{IV.58.7})$$

For $d = 4$ the same conclusion holds with $\mu_T = \mathbb{P}_z[\text{rank } \Omega_{T_z} \geq 4]$ and $\gamma_3 = 1/4$.

Proof. Bound each term in the exact slicing identity by its cap. For the second claim, the $2^{-r(T)}$ integrable fraction has cap one, the remaining fraction has cap b_{d-1} , and on the indicated subfraction the strict noncanonical cap subtracts at least γ_{d-1} . This is exactly (IV.58.7). $\square$

This profile can be iterated further or combined with the unsliced cap by taking their minimum. It is a valid finite certificate hierarchy, not a claim that independent optimal slice functions are mutually compatible or that the hierarchy computes the exact maximum.

IV.59 Finite calculus and the integrable tensor space

Source 69, delivered Section 2.

IV.59.1 Cube expansion, slicing, and phase twisting

Lemma IV.59.1 (Exact cube and slicing identities). *For a multilinear tensor and a complex function,*

$$\mathcal{E}_T(f) = \mathbb{E}_{x, h_1, \dots, h_d} \partial_{h_1} \cdots \partial_{h_d} f(x) (-1)^{T(h_1, \dots, h_d)}. \quad (\text{IV.59.1})$$

For symmetric T and $d \geq 3$,

$$\mathcal{E}_T(f) = \mathbb{E}_z \mathcal{E}_{T_z}(\partial_z f). \quad (\text{IV.59.2})$$

Proof. Fix $\mathbf{h} = (h_1, \dots, h_{d-1})$, and write $g = \partial_{h_1} \cdots \partial_{h_{d-1}} f$ and $\xi(x) = T(\mathbf{h}, x)$. Expand the squared Fourier coefficient and set one base point equal to the other plus y :

$$\left| \mathbb{E}_x g(x) (-1)^{\xi(x)} \right|^2 = \mathbb{E}_{x, y} g(x+y) \overline{g(x)} (-1)^{\xi(y)}.$$

Averaging over $\mathbf{h}$ proves (IV.59.1). Multiplicative derivatives commute, as follows by expanding their factors. Freeze a cube direction as z , move its derivative next to f , and use symmetry to put z in the contracted slot of T . Applying the cube identity in degree $d-1$ gives (IV.59.2). $\square$

No cancellation divides by f , so these identities hold at zero amplitudes. We use the normalization

$$\|f\|_{U^d}^{2^d} = \mathcal{E}_0(f). \quad (\text{IV.59.3})$$

Only this cube expression is needed below.

Lemma IV.59.2 (Gauge invariance). *If $S \in \mathcal{I}_d(V)$ has primitive P , then*

$$\mathcal{E}_{T+S}(fe(P)) = \mathcal{E}_T(f), \quad \mathcal{M}(T+S) = \mathcal{M}(T). \quad (\text{IV.59.4})$$

In particular, if T is integrable,

$$\mathcal{E}_T(f) = \|fe(-P_T)\|_{U^d}^{2^d}. \quad (\text{IV.59.5})$$

Proof. The d -fold multiplicative derivative of $e(P)$ equals $(-1)^S$. In (IV.59.1) its sign cancels the added S in the tensor. Multiplication by a unit-modulus phase permutes the unit polydisc of functions. This proves both identities in (IV.59.4), and the case with residual tensor zero gives (IV.59.5). $\square$

IV.59.2 The integration criterion, with a coordinate proof

Theorem IV.59.3 (Boolean nonclassical integration). *For $d \geq 3$, a symmetric tensor T is integrable if and only if $\Phi_T = 0$. Every symmetric bilinear tensor is integrable. In a chosen basis, the integrable tensors are exactly those whose basis coefficients depend only on the nonempty set of indices occurring, not on their positive multiplicities. The coefficient and primitive descriptions also apply at $d = 2$. Consequently*

$$\dim_{\mathbb{F}_2} \mathcal{I}_d(\mathbb{F}_2^n) = \sum_{j=1}^{\min(n, d)} \binom{n}{j}. \quad (\text{IV.59.6})$$

If c_A is the coefficient indexed by the support A , a primitive is

$$P_T(x) = \sum_{\substack{\emptyset \neq A \subseteq [n] \\ |A| \leq d}} \frac{c_A \prod_{i \in A} |x_i|}{2^{d-|A|+1}} \pmod{1}, \quad |x_i| \in \{0, 1\}. \quad (\text{IV.59.7})$$

This is the characteristic-two instance of the established nonclassical integration theorem [150]. The following proof also records the explicit basis used later in the sharp model construction.

Proof. Since translation by a has order two, additive difference operators on circle-valued functions satisfy

$$\Delta_a^2 = -2\Delta_a, \quad \Delta_a^2 \Delta_b = \Delta_a \Delta_b^2. \quad (\text{IV.59.8})$$

They commute. Every top derivative is therefore symmetric and obeys the repetition relation defining $\Phi_T = 0$.

Conversely, in a symmetric tensor with zero defect, the relation

$$T(e_i, e_i, e_j, \mathbf{z}) = T(e_i, e_j, e_j, \mathbf{z})$$

moves one occurrence from i to j whenever both occur and i has at least two occurrences. Any two positive multiplicity lists of a fixed total on a fixed support can be joined by such transfers. The coefficient is consequently a well-defined c_A . Conversely, support-dependent coefficients satisfy all these relations. For bilinear tensors the support already determines the multiplicity list, so no condition is needed.

It remains to check the primitive. For $a = |A|$, put $P_A(x) = 2^{-(d-a+1)} \prod_{i \in A} |x_i|$. Coordinate differences satisfy $\Delta_{e_i}^t = (-2)^{t-1} \Delta_{e_i}$ for $t \geq 1$, and

$$\Delta_h = \prod_{i: h_i=1} (I + \Delta_{e_i}) - I. \quad (\text{IV.59.9})$$

A term in an expansion of $d+1$ differences contains at least $d+1$ coordinate differences. A coordinate outside A kills the term. If the term uses $r \leq a$ distinct coordinates a total of $t \geq d+1$ times, its numerator has a factor 2^{t-r} and hence is divisible by the denominator 2^{d-a+1} . Thus it vanishes modulo one.

For exactly d coordinate differences, an index outside A again gives zero. If the support is a proper subset of A , the same divisibility argument gives an integer. If the support is exactly A , the value is $\pm 1/2 = 1/2 \pmod{1}$, independent of x . The top difference of a degree-at-most- d function is multilinear and base-point independent: the formula $\Delta_{a+b} = \Delta_a + \Delta_b + \Delta_a \Delta_b$ makes the additivity error a $(d+1)$ -fold difference. Hence agreement on basis tuples proves that (IV.59.7) has top derivative $T/2$. Finally, there is one freely chosen bit for each nonempty support of size at most d , proving (IV.59.6). $\square$

Lemma IV.59.4 (Contraction and diagonal lowering preserve integrability). *If $S \in \mathcal{I}_{m+1}(V)$ and $m \geq 2$, then both*

$$S_a(h_1, \dots, h_m) = S(a, h_1, \dots, h_m), \quad D(S)(h_1, \dots, h_m) = S(h_1, h_1, h_2, \dots, h_m) \quad (\text{IV.59.10})$$

belong to $\mathcal{I}_m(V)$.

Proof. If P is a primitive of S , then $\Delta_a P$ is a degree-at-most- m primitive of S_a ; equivalently, the repetition relation persists on freezing one slot.

For $D(S)$, the two cross terms in the repeated slot cancel, so the map is multilinear. The repetition relation for S moves the doubled variable between slots, proving symmetry of $D(S)$. On a basis tuple, duplicating one index does not change its support. Therefore the coefficients of $D(S)$ depend only on support, and Theorem IV.59.3 gives integrability. In fact D retains exactly the support coefficients of sizes at most m and discards those of size $m + 1$. $\square$

The diagonal-lowering observation will control the *dimension* of a space of defects, rather than merely provide another integration test.

IV.60 The defect and the cubic spectral base case

Source 69, delivered Section 3.

IV.60.1 Alternation and compatibility

Lemma IV.60.1 (Structure of the repeated defect). *For symmetric T , the map $\Phi_T : V^{d-3} \rightarrow \text{Alt}^2(V)$ is symmetric multilinear. For $d \geq 4$ it satisfies*

$$\Phi_T(c, \mathbf{w})(a, b) + \Phi_T(b, \mathbf{w})(a, c) + \Phi_T(a, \mathbf{w})(b, c) = 0, \quad (\text{IV.60.1})$$

where $\mathbf{w} \in V^{d-4}$.

Proof. In $T(a + a', a + a', b, z)$ the two cross terms agree by symmetry and cancel in $\mathbb{F}_2$. The same holds for the repeated b term. The defect is bilinear in a, b , is symmetric in that pair, and vanishes when $a = b$; hence it is alternating. Linearity and symmetry in the frozen variables are inherited from T . Expanding (IV.60.1) produces six tensor values. Symmetry pairs them into three equal pairs, so the sum is zero. $\square$

IV.60.2 A projective orthogonality estimate

The next argument is the finite-dimensional spectral method used in the preceding project manuscripts [283, 284]. We include it to make the cubic base case independent of an unreviewed reference.

For a bilinear $B : V^2 \rightarrow \mathbb{F}_2$, put $\Omega_B(a, b) = B(a, b) + B(b, a)$ and define

$$Q_B(g, h) = \mathbb{E}_a \left| \mathbb{E}_x g(x + a) \overline{h(x)} (-1)^{B(a, x)} \right|^2, \quad Q_B(g) = Q_B(g, g). \quad (\text{IV.60.2})$$

Lemma IV.60.2 (Projective orthogonality). *If the alternating form Ω_B has rank $2r$, then*

$$Q_B(g, h) \leq 2^{-r} \|g\|_2^2 \|h\|_2^2. \quad (\text{IV.60.3})$$

Also

$$Q_B(g) = \mathbb{E}_{x, a, b} \partial_a \partial_b g(x) (-1)^{B(a, b)}. \quad (\text{IV.60.4})$$

Proof. Use the normalized Hilbert space with inner product $\langle g, h \rangle = \mathbb{E}_x g(x) \overline{h(x)}$, linear in its first argument. Define

$$(U_a g)(x) = (-1)^{B(a,x)} g(x+a).$$

These operators are unitary and satisfy

$$U_a U_b = (-1)^{B(b,a)} U_{a+b}, \quad U_a U_b = (-1)^{\Omega_B(a,b)} U_b U_a.$$

The positive operator

$$M_g v = \mathbb{E}_a \langle v, U_a g \rangle U_a g$$

has trace $\|g\|_2^2$. It commutes with every U_b , since conjugation by U_b permutes its rank-one summands and the scalar signs cancel.

An alternating form splits into r nondegenerate two-dimensional pairs and its radical. This follows by choosing a, b with pairing one, taking the orthogonal complement of their span, and inducting. In particular the rank is even. Choose a subspace H containing the radical and one isotropic line from each pair. Then $H = H^\perp$ and $[V : H] = 2^r$. Every nonzero common invariant subspace for the U_a has dimension at least 2^r : the commuting unitaries U_h , $h \in H$, have a common eigenvector, and its translates by representatives of V/H have distinct joint eigenvalues. Their eigenvalue ratios on h are $(-1)^{\Omega_B(h,a)}$. Distinct cosets give distinct ratios, hence orthogonal vectors.

Every positive eigenspace of M_g is such an invariant subspace. Therefore each positive eigenvalue has multiplicity at least 2^r , and $\|M_g\|_{\text{op}} \leq 2^{-r} \text{tr } M_g$. Since $Q_B(g, h) = \langle M_g h, h \rangle$, this proves (IV.60.3). Expanding the square in (IV.60.2) and replacing its second base point by $x + b$ gives the complex conjugate of the right side of (IV.60.4). It is real because it equals $Q_B(g)$, so the conjugation can be removed. $\square$

Theorem IV.60.3 (Exact cubic obstruction law). *For $T \in \text{SM}_3(V)$ let $\Omega_T(a, b) = T(a, a, b) + T(a, b, b)$ have rank $2r$. Then*

$$\mathcal{M}(T) = 2^{-r}. \quad (\text{IV.60.5})$$

In particular, a nonintegrable symmetric cubic has energy at most $1/2$.

Proof. Write $\tau_w f(x) = f(x + w)$. The product identity

$$\partial_a \partial_{a+w} f = \overline{\partial_a (f \tau_w f)} \quad (\text{IV.60.6})$$

is valid even when f vanishes. In the cubic expression (IV.59.1), put the third direction equal to $a + w$. The sign is then given by

$$B_w(a, b) = T(a, a, b) + T(a, b, w).$$

This is bilinear, and its alternating part is Ω_T because the term containing w is symmetric in a, b . Identity (IV.60.6) and Lemma IV.60.2 give

$$\mathcal{E}_T(f) = \mathbb{E}_w Q_{B_w}(f \tau_w f) \leq 2^{-r} \mathbb{E}_w \|f \tau_w f\|_2^4 \leq 2^{-r}.$$

For equality, split Ω_T into independent symplectic pairs $\sum_{i=1}^r (u_i \otimes v_i + v_i \otimes u_i)$. The tensor $\sum_i C_3^{u_i, v_i}$ has exactly this defect, as one checks by substitution. Their difference is integrable by Theorem IV.59.3. For the constant function, the cube average for the sum of these tensors factors over the independent pairs. On one pair its value is $1/2$: the selected frequency is nonzero on half of the pairs of directions. Thus the product has energy 2^{-r} . Twisting by a primitive of the integrable difference gives equality for T . This use of a product gives a lower bound by an explicit function; it makes no assertion that arbitrary optimizing functions factor. $\square$

IV.61 Rank-one defects and canonical energies

Source 69, delivered Section 4.

IV.61.1 The contraction radical and its rank-one case

For $d \geq 4$, Theorem IV.59.3 gives

$$R(T) = \{z : \Phi_T(z, z_2, \dots, z_{d-3}) = 0 \text{ for every } z_2, \dots, z_{d-3}\}. \quad (\text{IV.61.1})$$

It is a linear subspace. By symmetry it is a radical in every frozen slot. It equals V precisely when T is integrable. Hence a nonintegrable tensor has $\rho(T) \geq 1$.

Lemma IV.61.1 (A one-form compatibility identity). *If $u \neq 0$ and $\Omega \in \text{Alt}^2(V)$ satisfy*

$$u(c)\Omega(a, b) + u(b)\Omega(a, c) + u(a)\Omega(b, c) = 0 \quad \text{for every } a, b, c,$$

then $\Omega = u \otimes v + v \otimes u$ for some $v \in V^$. If $\Omega \neq 0$, u, v are independent and Ω has rank two.*

Proof. Choose c_0 with $u(c_0) = 1$ and put $v(x) = \Omega(x, c_0)$. Substitution of c_0 gives the claimed identity. The nonzero form has a two-dimensional nondegenerate quotient, namely that detected by u, v . $\square$

Proposition IV.61.2 (Rank-one classification; preceding project result). *For $T \in \text{SM}_d(V) \setminus \mathcal{I}_d(V)$, $d \geq 4$, the following are equivalent: $\rho(T) = 1$; the defect has the form*

$$\Phi_T(\mathbf{z}) = \Omega \prod_i u(z_i), \quad \Omega \neq 0, \quad u \neq 0; \quad (\text{IV.61.2})$$

and $T = C_d^{u,v} + S$ for independent u, v and integrable S .

Proof. If the frozen radical has codimension one, multilinearity makes the defect descend in all frozen slots to the one-dimensional quotient. Hence it has form (IV.61.2). Conversely that form has frozen radical $\ker u$. In (IV.60.1), set the other frozen variables to vectors on which u equals one. Lemma IV.61.1 gives $\Omega = u \otimes v + v \otimes u$. A direct expansion of (F.74.5) gives

$$\Phi_{C_d^{u,v}}(\mathbf{z})(a, b) = \left(\prod_i u(z_i) \right) (u(a)v(b) + v(a)u(b)). \quad (\text{IV.61.3})$$

Indeed the two terms placing v in repeated a slots cancel, as do the two copies of every term placing v in a frozen slot. The surviving terms are exactly the displayed ones. Thus $T - C_d^{u,v}$ has zero defect and is integrable. Formula (IV.61.3) also proves the converse. $\square$

This proposition was proved in [284]. The improvement below comes from how the higher-rank alternative is handled.

IV.61.2 Periods of derivatives

Lemma IV.61.3 (Exact even-direction periods). *Let $F = \partial_{h_1} \cdots \partial_{h_t} f$, $t \geq 1$. Then*

$$F(x + h_i) = \overline{F(x)}, \quad F(x + w) = F(x) \quad \text{if } w \in \text{span}\{h_1 + h_i : 2 \leq i \leq t\}. \quad (\text{IV.61.4})$$

Consequently a Fourier coefficient of F at $\xi \in V^$ can be nonzero only if $\xi(h_1) = \cdots = \xi(h_t)$.*

Proof. The cube product defining F has a factor at each labelled vertex, with conjugation parity given by the vertex parity and t . Translation by h_i toggles vertex coordinate i and reverses every conjugation parity. This conjugates the product. Two such translations leave it unchanged, which proves the period assertion, including when vertices coincide or amplitudes vanish. Translating a Fourier sum by a period w with $\xi(w) = 1$ changes its sign without changing F , so that coefficient is zero. $\square$

Lemma IV.61.4 (A pure tensor on a constrained derivative). *Let $U_m(h_1, \dots, h_m) = \prod_i u(h_i)$ with $u \neq 0$ and $m \geq 2$. If $z \in \ker u$ and f is one-bounded, then*

$$\mathcal{E}_{U_m}(\partial_z f) \leq 1 - 2^{1-m}. \quad (\text{IV.61.5})$$

The constant function attains this value.

Proof. The selected frequency for U_m is u exactly when all $m - 1$ directions h_i have $u(h_i) = 1$; otherwise it is zero. On the first event the function being transformed is $\partial_{h_1} \cdots \partial_{h_{m-1}} \partial_z f$. It is invariant under $h_1 + z$, but $u(h_1 + z) = 1$. Its selected coefficient therefore vanishes by Lemma IV.61.3. The event has probability 2^{1-m} , and every remaining squared coefficient is at most one. For $f = 1$, zero-frequency coefficients are one and nonzero-frequency coefficients are zero, proving equality. $\square$

Without the constraint that the tested function is $\partial_z f$, the maximum for U_m is one: it has primitive $|u(x)|/2^m$. This distinction is essential to the canonical extremum and was the key advance of [284].

IV.61.3 The dimension-free canonical maximum

Theorem IV.61.5 (Canonical maximum; preceding project result). *For independent u, v , every $d \geq 3$, and any ambient dimension,*

$$\mathcal{M}(C_d^{u,v}) = \mathcal{E}_{C_d^{u,v}}(1) = b_d. \quad (\text{IV.61.6})$$

Proof. For the constant function, the selected frequency has the form $Av + Bu$, where, with $t = d - 1$,

$$A = \prod_{j=1}^t u(h_j), \quad B = \sum_{i=1}^t v(h_i) \prod_{j \neq i} u(h_j).$$

If at least two $u(h_j)$ vanish, both coefficients vanish. If exactly one vanishes, $A = 0$ and B is an independent unbiased bit. If none vanishes, $A = 1$. Thus the probability of a nonzero frequency is $2^{-t} + t2^{-t-1} = (d+1)/2^d$, giving the claimed lower bound.

The cubic upper bound follows from Theorem IV.60.3 and the rank-two defect in (IV.61.3). Suppose the upper bound is known in degree $m = d - 1 \geq 3$. The exact tensor contraction is

$$(C_d^{u,v})_z = u(z)C_m^{u,v} + v(z)U_m.$$

The four pairs $(u(z), v(z))$ occur equally often. On $(0, 0)$ the slice energy is at most one. On $(0, 1)$ it is at most $1 - 2^{1-m}$ by Lemma IV.61.4. On the half with $u(z) = 1$, remove the integrable tensor $v(z)U_m$ by gauge invariance and use the inductive bound b_m . Slicing therefore gives

$$\begin{aligned}\mathcal{E}_{C_d^{u,v}}(f) &\leq \frac{1}{4} + \frac{1}{4}(1 - 2^{2-d}) + \frac{1}{2}b_{d-1} \\ &= 1 - \frac{d+1}{2^d}.\end{aligned}\tag{IV.61.7}$$

This proof allows f to depend on every ambient coordinate. It never reduces the optimization problem to a two-dimensional space. $\square$

IV.62 The contraction-integrability induction

Source 69, delivered Section 5.

IV.62.1 A general one-step inequality

Lemma IV.62.1 (Counting nonintegrable contractions). *Fix $d \geq 4$, and suppose that every nonintegrable symmetric tensor of degree $d-1$ on V has maximum energy at most β . For a symmetric tensor T with $\rho(T) = r$,*

$$\mathcal{E}_T(f) \leq 2^{-r} + (1 - 2^{-r})\beta.\tag{IV.62.1}$$

Proof. The subspace $R(T)$ has proportion 2^{-r} in V . Its contractions are integrable, so the trivial bound one suffices there. Every contraction outside it is nonintegrable and has energy at most β when tested on $\partial_z f$, which is still one-bounded. Average these two bounds in (IV.59.2). $\square$

The test of whether T_z is integrable is an *identity* on all its remaining variables. Replacing it by the probability that a randomly completed defect is nonzero would give a different, weaker estimate.

IV.62.2 Proof of the all-degree theorem

Proof of Theorem F.74.1. Theorem IV.60.3 gives the universal bound $b_3 = 1/2$. Proceed by induction on $d \geq 4$. Let T be nonintegrable. If $\rho(T) = 1$, Proposition IV.61.2 and Theorem IV.61.5, followed by gauge invariance, give $\mathcal{M}(T) = b_d$.

If $\rho(T) \geq 2$, apply Lemma IV.62.1 with $\beta = b_{d-1} = 1 - d/2^{d-1}$. It gives the more precise estimate

$$\boxed{\mathcal{E}_T(f) \leq 1 - \frac{d}{2^{d-1}}(1 - 2^{-\rho(T)})}.\tag{IV.62.2}$$

Since $1 - 2^{-\rho(T)} \geq 3/4$,

$$\mathcal{E}_T(f) \leq 1 - \frac{3d}{2^{d+1}} = b_d - \frac{d-2}{2^{d+1}} < b_d.$$

This proves the inductive universal bound and excludes every higher-rank tensor from equality. The canonical example attains b_d and is nonintegrable by (IV.61.3). It proves sharpness in every $n \geq 2$ and completes the classification. $\square$

Remark IV.62.2 (Why a long support census is unnecessary here). The preceding proof of (F.74.15) uses a support bound for non-pure symmetric multilinear defects. That bound treats a nonzero contraction by how rarely its *values* can be nonzero. In contrast, once the lower-degree energy theorem is available, every nonintegrable contraction already costs a full loss $1 - b_{d-1}$. Unless the frozen radical is a hyperplane, at least three quarters of all contractions are of this type. The hyperplane case is exactly the canonical family, for which the separate derivative-period argument is sharp. These two alternatives close the induction without a new tensor-support estimate.

IV.62.3 A rank-sensitive statement and a one-step deficit identity

Estimate (IV.62.2) is valid even when $\rho = 1$, although it is then weaker than the exact canonical bound. At a fixed degree it improves the rank-one alternative for every $\rho \geq 2$. For $r = \rho(T)$, write

$$c_d(r) = 2^{-r} + (1 - 2^{-r})b_{d-1}.$$

The proof supplies an exact nonnegative decomposition:

$$\begin{aligned} c_d(r) - \mathcal{E}_T(f) &= \mathbb{E}_z \mathbf{1}_{R(T)}(z)(1 - \mathcal{E}_{T_z}(\partial_z f)) \\ &\quad + \mathbb{E}_z \mathbf{1}_{V \setminus R(T)}(z)(b_{d-1} - \mathcal{E}_{T_z}(\partial_z f)). \end{aligned} \quad (\text{IV.62.3})$$

Both integrands are nonnegative. Thus, for example, if $\mathcal{E}_T(f) \geq c_d(r) - \delta$, the fraction of all directions outside $R(T)$ whose slice energy is less than $b_{d-1} - t$ is at most δ/t . This is a direct finite averaging consequence, not a phase-stability theorem.

Corollary IV.62.3 (Exact phase-removal certificate). *If $T \in \text{SM}_d(V)$ and $\mathcal{E}_T(f) > b_d$, then $T \in \mathcal{I}_d(V)$ and the explicit primitive in (IV.59.7) satisfies*

$$\|fe(-P_T)\|_{U^d}^{2d} = \mathcal{E}_T(f).$$

For $d \geq 3$ and $n \geq 2$, the strict threshold cannot be lowered.

Proof. Theorem F.74.1 excludes nonintegrability, and gauge invariance then gives equality. At the threshold $C_d^{u,v}$ and $f = 1$ remain nonintegrable, proving optimality. $\square$

IV.63 Extremal flags and quantitative tensor rigidity

Source 69, delivered Section 6.

IV.63.1 Flags, rather than merely two-dimensional quotients

Two tensors are *gauge-equivalent* when their difference belongs to $\mathcal{I}_d(V)$. By Theorem IV.59.3, this is equivalent to having identical defects.

Theorem IV.63.1 (Enumeration of extremal gauge classes). *For $d \geq 4$, the gauge classes attaining the nonintegrable maximum b_d are in bijection with flags*

$$\ell \subset \Pi \subset V^*, \quad \dim \ell = 1, \quad \dim \Pi = 2. \quad (\text{IV.63.1})$$

Their number is

$$(2^n - 1)(2^{n-1} - 1). \quad (\text{IV.63.2})$$

The class represented by $C_d^{u,v}$ corresponds to $\ell = \text{span}(u)$ and $\Pi = \text{span}(u, v)$.

Proof. The radical of a canonical defect is $\ker u$, so the defect determines the nonzero functional u uniquely over $\mathbb{F}_2$. Choose w with $u(w) = 1$. Evaluating every frozen slot at w recovers the alternating form $\Omega = u \otimes v + v \otimes u$. For fixed u , this determines v modulo $\text{span}(u)$: if $u \wedge (v + v') = 0$, choose c with $u(c) = 1$ and evaluate at (c, x) to get $v + v' = (v(c) + v'(c))u$. Thus the defect determines exactly the stated flag. Conversely a flag specifies u and v modulo u , and formula (IV.61.3) then specifies the same defect. Theorem F.74.1 proves that these are all extremal classes. There are $2^n - 1$ choices of $u \neq 0$ and $2^{n-1} - 1$ nonzero cosets of $\text{span}(u)$ in V^* . This proves (IV.63.2). $\square$

For $d = 3$ there is no frozen slot to distinguish a line inside the plane. The maximal nonintegrable cubic classes are the rank-two alternating forms, indexed by planes in V^* . Each plane has three lines, so the cubic count is $(2^n - 1)(2^{n-1} - 1)/3$. Confusing these two classifications would overcount the cubic endpoint by a factor of three.

IV.63.2 A discrete rigidity band

Corollary IV.63.2 (High-energy tensor dichotomy). *Let $d \geq 4$. If*

$$\mathcal{E}_T(f) > 1 - \frac{3d}{2^{d+1}}, \quad (\text{IV.63.3})$$

then either T is integrable or it is gauge-equivalent to exactly one extremal flag class. In particular, a nonintegrable tensor whose energy lies at distance strictly less than $(d - 2)/2^{d+1}$ from b_d must already have the exact canonical defect.

Proof. A nonintegrable tensor outside the canonical classes satisfies (F.74.8). Strict inequality excludes that alternative; uniqueness of the flag follows from Theorem IV.63.1. $\square$

This is rigidity of a discrete tensor invariant. It does not assert that the function f is close to a particular phase. Such a conclusion would require an analysis of equality and stability inside the canonical spectral problem.

IV.63.3 A small exact example

On $V = \mathbb{F}_2^3$, symmetric quartic tensors have $\binom{6}{4} = 15$ free coefficients. The integrable subspace has dimension $3 + 3 + 1 = 7$. There are 21 extremal gauge classes and therefore $21 \cdot 2^7 = 2688$ rank-one-defect tensors. The accompanying exhaustive coefficient calculation gives

$\rho(T)$	Number of tensors	$m_{\text{eff}}(T)$
0	128	0
1	2688	2
2	19712	3
3	10240	3

The first two counts also follow from the preceding proofs. The last two are reported as exact finite computations, not used in any general theorem.

IV.64 The minimal obstruction quotient

Source 69, delivered Section 7.

IV.64.1 The complete defect radical

The frozen radical $R(T)$ need not annihilate the two alternating slots. A direction can be invisible when used as a frozen variable but still enter the obstruction through one of those slots. This is why $\rho(T)$ alone need not equal the minimal model dimension. The subspace $K(T)$ in (F.74.13) imposes both requirements. Symmetry of the frozen slots and alternation of the last pair make it precisely the common radical of the defect in *every* slot.

Theorem IV.64.1 (Exact minimal quotient). *For $d \geq 4$,*

$$m_{\text{eff}}(T) = \text{codim}_V K(T). \quad (\text{IV.64.1})$$

*Moreover the quotient $q : V \rightarrow V/K(T)$ itself supports a model $T = S + q^*T_0$ with S integrable.*

Proof. If $T = S + q^*T_0$ and S is integrable, then $\Phi_T = q^*\Phi_{T_0}$. Every vector in $\ker q$ kills the defect in all slots. Thus $\ker q \subseteq K(T)$ and the target dimension is at least $\text{codim } K(T)$.

Conversely, multilinearity and the definition of $K(T)$ show that the defect descends in every slot to $V/K(T)$. Choose a linear section $j : V/K(T) \rightarrow V$ of the quotient map and define $T_0 = j^*T$. The defect of q^*T_0 equals Φ_T , since every vector x and $j q(x)$ differ by an element of the complete radical. Therefore $S = T - q^*T_0$ has zero defect and is integrable. This produces a model in exactly the lower-bound dimension. $\square$

A different section may change the representative T_0 , but not the minimal dimension or the defect on the quotient. Importantly, the theorem is about representing the *tensor modulo integrable tensors*. It does not imply that optimizing functions descend to this quotient, nor that $\mathcal{M}(q^*T_0) = \mathcal{M}(T_0)$ for arbitrary T_0 .

IV.64.2 Two radical arguments force zero

Lemma IV.64.2 (The defect vanishes on two frozen-radical directions). *If $a, b \in R(T)$, then $\Phi_T(z)(a, b) = 0$ for every frozen tuple.*

Proof. Take the first frozen variable to be c and the remaining ones to be $\mathbf{w}$. Equation (IV.60.1) expresses $\Phi_T(c, \mathbf{w})(a, b)$ as the sum of two terms having either a or b in a frozen slot. Both terms vanish by (IV.61.1). $\square$

IV.64.3 The integrable-valued radical map

Choose a complement H with $V = H \oplus R(T)$ and $\dim H = \rho(T)$. For $a \in R(T)$ define a map on H^{d-2} by

$$\Psi_a(h_1, \dots, h_{d-2}) = \Phi_T(h_2, \dots, h_{d-2})(a, h_1). \quad (\text{IV.64.2})$$

Proposition IV.64.3 (An integrable-valued linear map). *The assignment*

$$\Psi : R(T) \longrightarrow \mathcal{I}_{d-2}(H), \quad a \longmapsto \Psi_a \quad (\text{IV.64.3})$$

is linear and has kernel $K(T)$.

Proof. Linearity in a follows from bilinearity of the alternating pair in the defect. Since $a \in R(T)$, the tensor $S_a = T_a$ is integrable of degree $d - 1$. Expanding the definition gives, with $\mathbf{z} = (h_2, \dots, h_{d-2})$,

$$\Psi_a(h_1, \mathbf{z}) = S_a(a, h_1, \mathbf{z}) + S_a(h_1, h_1, \mathbf{z}). \quad (\text{IV.64.4})$$

The first summand is an ordinary contraction of S_a , and the second is its diagonal lowering. Both are integrable by Lemma IV.59.4; restriction to H preserves that property. Thus (IV.64.3) has exactly the indicated, smaller target space.

If $a \in K(T)$, then certainly $\Psi_a = 0$. Conversely suppose $a \in R(T)$ and $\Psi_a = 0$. Every frozen variable may be reduced to its H component, because its $R(T)$ component is in a frozen radical. The other alternating variable may also be reduced to H , since its $R(T)$ component pairs with $a \in R(T)$ and vanishes by Lemma IV.64.2. The remaining values are exactly those of Ψ_a , hence all zero. Therefore $a \in K(T)$. $\square$

The assertion that Ψ_a is integrable, not merely symmetric multilinear, is the point that saves coordinates. Bounding its target by all symmetric $(d - 2)$ -linear tensors would give $\rho + \binom{\rho+d-3}{d-2}$ instead of (F.74.14).

Proof of the upper bound in Theorem F.74.2. Proposition IV.64.3 and rank-nullity give the identity

$$m_{\text{eff}}(T) = \text{codim } K(T) = \rho(T) + \text{rank}(\Psi : R(T) \rightarrow \mathcal{I}_{d-2}(H)). \quad (\text{IV.64.5})$$

Apply the dimension formula (IV.59.6) to the target. This proves (F.74.14), including the integrable case $\rho = 0$. $\square$

IV.64.4 A complete normal form for a prescribed radical

The preceding map is part of a more precise description of the quotient by integrable tensors. Here the radical is only *prescribed to contain* a fixed subspace; it need not be exactly that subspace.

Theorem IV.64.4 (Fixed-radical normal form). *Fix a splitting $V = H \oplus R_0$, put $r = \dim H$ and $t = \dim R_0$, and let*

$$\mathcal{W}_d(H, R_0) = \{T \in \text{SM}_d(V) : R_0 \subseteq R(T)\}, \quad d \geq 4.$$

For $a \in R_0$ define Ψ_a by (IV.64.2), using H as written. There is a linear isomorphism

$$\mathcal{W}_d(H, R_0)/\mathcal{I}_d(V) \cong (\text{SM}_d(H)/\mathcal{I}_d(H)) \oplus \text{Hom}_{\mathbb{F}_2}(R_0, \mathcal{I}_{d-2}(H)), \quad (\text{IV.64.6})$$

given by $[T] \mapsto ([T|_H], \Psi)$. In particular its dimension is

$$\binom{r+d-1}{d} - \sum_{j=1}^{\min(r,d)} \binom{r}{j} + t \sum_{j=1}^{\min(r,d-2)} \binom{r}{j}. \quad (\text{IV.64.7})$$

Proof. The space $\mathcal{I}_d(V)$ is contained in $\mathcal{W}_d(H, R_0)$, and both restriction and Ψ respect addition of an integrable tensor. The proof of Proposition IV.64.3 shows that Ψ_a is integrable for every $a \in R_0$; only the inclusion $R_0 \subseteq R(T)$ is needed. Thus the displayed map is well defined and linear.

If its image is zero, the defect vanishes with all arguments in H , and also when one alternating argument belongs to R_0 and the other to H . It vanishes when both alternating arguments belong to R_0 by the Bianchi identity, and whenever a frozen argument belongs to R_0 by hypothesis. Expand every argument into its H and R_0 components. These cases cover all terms, proving $\Phi_T = 0$. Hence the map is injective.

For surjectivity, choose a representative $A \in \text{SM}_d(H)$ of the desired first component and a map $L : R_0 \rightarrow \mathcal{I}_{d-2}(H)$ for the second component. The diagonal-lowering map $D : \mathcal{I}_{d-1}(H) \rightarrow \mathcal{I}_{d-2}(H)$ is surjective: retain the prescribed support coefficients of sizes at most $d-2$ and set the support coefficients of size $d-1$ equal to zero. Choose a basis $r_1, \dots, r_t$ of R_0 and lifts $S_i \in \mathcal{I}_{d-1}(H)$ with $D(S_i) = L(r_i)$. Define a symmetric tensor B to be zero on basis tuples with zero or at least two R_0 slots, and set

$$B(r_i, h_1, \dots, h_{d-1}) = S_i(h_1, \dots, h_{d-1})$$

on the remaining basis tuples. The contractions B_{r_i} are pullbacks of the integrable tensors S_i , so $R_0 \subseteq R(B)$. Moreover $B|_{H^d} = 0$, and the term containing two R_0 slots vanishes in (IV.64.4). Thus $\Psi_{r_i}^B = D(S_i) = L(r_i)$. Adding the pullback of A along $V \rightarrow H$ produces exactly the desired pair of components. This proves surjectivity. The dimensions of the two factors, using (IV.59.6), give (IV.64.7). $\square$

The theorem gives a full linear parameterization before imposing changes of coordinates. Requiring $R(T) = R_0$, rather than just containment, is an additional nondegeneracy condition on the pair of parameters. The sharpness construction below enforces this condition with singleton support tensors.

IV.65 A construction attaining every dimension bound

Source 69, delivered Section 8.

IV.65.1 Support-basis construction

Fix $d \geq 4$ and $r \geq 1$, put $m = d - 2$, and let $H = \mathbb{F}_2^r$ with basis $e_1, \dots, e_r$. Index the family

$$\mathcal{A} = \{A \subseteq [r] : 1 \leq |A| \leq m\}, \quad L = |\mathcal{A}| = \sum_{j=1}^{\min(r,m)} \binom{r}{j}. \quad (\text{IV.65.1})$$

For each $A \in \mathcal{A}$ let $S_A \in \mathcal{I}_{d-1}(H)$ be the support-basis tensor whose coefficient on a $(d-1)$ -tuple of basis vectors is one exactly when the set of its indices is A . Let R_0 have basis $\{r_A : A \in \mathcal{A}\}$, and take $V = H \oplus R_0$. Define $T \in \text{SM}_d(V)$ by its unordered basis coefficients as follows:

$$T(r_A, e_{i_1}, \dots, e_{i_{d-1}}) = S_A(e_{i_1}, \dots, e_{i_{d-1}}); \quad T = 0 \text{ on tuples with zero or at least two } R_0 \text{ slots.} \quad (\text{IV.65.2})$$

Extend symmetrically and multilinearly. This is a complete explicit specification of a tensor on $r + L$ coordinates.

Theorem IV.65.1 (Sharpness of the effective-dimension bound). *The tensor in (IV.65.2) satisfies*

$$R(T) = R_0, \quad \rho(T) = r, \quad K(T) = 0, \quad m_{\text{eff}}(T) = r + L = B_d(r). \quad (\text{IV.65.3})$$

Proof. For each A , the contraction T_{r_A} is the pullback of S_A from H . It is integrable, so $R_0 \subseteq R(T)$. Take a nonzero $h \in H$ and choose i whose coordinate in h is one. Singletons belong to $\mathcal{A}$. From the definition, with all the remaining $d - 4$ frozen directions equal to e_i ,

$$\Phi_T(h, e_i, \dots, e_i)(r_{\{i\}}, e_i) = 1. \quad (\text{IV.65.4})$$

Indeed the term with two occurrences of $r_{\{i\}}$ is zero, and the other is the singleton support tensor evaluated on h and $d - 2$ copies of e_i . It equals the i th coordinate of h . Thus no nonzero H component lies in the frozen radical, and $R(T) = R_0$.

For $a = r_A$ the first summand in (IV.64.4) is zero, because it contains two R_0 basis vectors. Hence

$$\Psi_{r_A}(h_1, \dots, h_m) = S_A(h_1, h_1, h_2, \dots, h_m) = D(S_A)(h_1, \dots, h_m).$$

By the support-coefficient description, the $D(S_A)$, $A \in \mathcal{A}$, are exactly the support basis of $\mathcal{I}_m(H)$. The map Ψ is therefore an isomorphism $R_0 \rightarrow \mathcal{I}_m(H)$. Proposition IV.64.3 gives $K(T) = 0$. The minimal quotient theorem yields $m_{\text{eff}}(T) = \dim V = r + L$. $\square$

This completes the proof of Theorem F.74.2. Notice that sharpness is algebraic: it does not require these examples to have unusually large Fourier energy.

IV.65.2 An explicit five-coordinate quartic

For $d = 4$ and $r = 2$, write the basis as e_1, e_2, a, b, c , corresponding to $\{1\}, \{2\}, \{1, 2\}$. The only nonzero unordered coefficients are

$$T(a, e_1, e_1, e_1) = T(b, e_2, e_2, e_2) = T(c, e_1, e_1, e_2) = T(c, e_1, e_2, e_2) = 1. \quad (\text{IV.65.5})$$

Here $R(T) = \text{span}(a, b, c)$ and $\rho(T) = 2$. The map Ψ sends these three basis vectors to the symmetric bilinear forms with coefficient patterns x_1y_1 , x_2y_2 , and $x_1y_2 + x_2y_1$, respectively. Thus $K(T) = 0$: no model modulo an integrable tensor can use fewer than all five coordinates. A two-dimensional frozen quotient can therefore hide three additional active directions in the alternating pair.

IV.65.3 Consequences for the size of models

If $d \geq r + 2$, then all nonempty subsets occur and

$$B_d(r) = r + 2^r - 1. \quad (\text{IV.65.6})$$

The construction proves that no bound polynomial in r , uniformly in the degree, can hold for the number of active coordinates modulo integrable tensors under this hypothesis alone. For a fixed $d \geq 4$, however,

$$B_d(r) = \frac{r^{d-2}}{(d-2)!} + O_d(r^{d-3} + r), \quad r \rightarrow \infty.$$

The two regimes concern different quantifiers and should not be conflated. Some small values are

$d \backslash r$	1	2	3	4	5
4	2	5	9	14	20
5	2	5	10	18	30
6	2	5	10	19	35
7	2	5	10	19	36

Each entry is attained in that ambient dimension. The case $r = 1$ recovers the two-coordinate canonical model, whereas the higher-rank examples show why simply extending that two-coordinate conclusion is impossible.

IV.66 Energy-to-complexity and robust selector certificates

Source 69, delivered Section 9.

IV.66.1 From energy to a finite obstruction model

Corollary IV.66.1 (Energy bounds the contraction rank). *Let $d \geq 4$, let $q \geq 0$ be an integer, and let $T \in \text{SM}_d(V)$. If*

$$\mathcal{E}_T(f) > 1 - \frac{d}{2^{d-1}}(1 - 2^{-(q+1)}), \quad (\text{IV.66.1})$$

then $\rho(T) \leq q$ and $m_{\text{eff}}(T) \leq B_d(q)$. Alternatively, if $\varepsilon = 1 - \mathcal{E}_T(f)$ and $2^{d-1}\varepsilon/d < 1$, then

$$\rho(T) \leq \log_2 \left(\frac{1}{1 - 2^{d-1}\varepsilon/d} \right). \quad (\text{IV.66.2})$$

Proof. If T is integrable, $\rho = m_{\text{eff}} = 0$. Otherwise (IV.62.2) applies. Its upper bound decreases with ρ , so (IV.66.1) rules out $\rho \geq q+1$. The model bound follows from Theorem F.74.2 and monotonicity of B_d . Rearranging the same inequality gives $2^{-\rho} \geq 1 - 2^{d-1}\varepsilon/d$ and hence (IV.66.2). $\square$

For $q = 0$ the exact integrability certificate at b_d is stronger. For $q = 1$ the conclusion can be sharpened to the integrable-or-canonical dichotomy rather than just a two-coordinate model. Formula (IV.66.2) is useful only where its right-hand side is finite; it is not a small-energy inverse theorem.

IV.66.2 Corrupted frequency selectors

For any selector $\sigma : V^{d-1} \rightarrow V^*$, define

$$\mathcal{E}_\sigma(f) = \mathbb{E}_{\mathbf{h}} \left| \mathbb{E}_x \partial_{\mathbf{h}} f(x) (-1)^{\sigma(\mathbf{h})(x)} \right|^2.$$

Here $\partial_{\mathbf{h}} = \partial_{h_1} \cdots \partial_{h_{d-1}}$. Let $\sigma_T(\mathbf{h}) = T(\mathbf{h}, \cdot)$.

Proposition IV.66.2 (Exact corruption margin). *If $\mathbb{P}_{\mathbf{h}}[\sigma(\mathbf{h}) \neq \sigma_T(\mathbf{h})] \leq \eta$, then*

$$|\mathcal{E}_\sigma(f) - \mathcal{E}_T(f)| \leq \eta. \quad (\text{IV.66.3})$$

Thus, for symmetric T , energy $\mathcal{E}_\sigma(f) > b_d + \eta$ certifies integrability and yields a primitive with

$$\|fe(-P_T)\|_{U^d}^{2^d} \geq \mathcal{E}_\sigma(f) - \eta. \quad (\text{IV.66.4})$$

For $d \geq 4$, energy $\mathcal{E}_\sigma(f) > 1 - 3d/2^{d+1} + \eta$ certifies that T is integrable or canonical modulo an integrable tensor.

Proof. Each squared coefficient lies in $[0, 1]$. The two averages agree outside the disagreement set, and their summands differ by at most one inside it. This proves (IV.66.3); the previous theorems give the remaining claims. $\square$

The disagreement event compares *whole linear functionals* at a tuple of derivative directions. A different hypothesis measuring disagreements of the scalar values $T(\mathbf{h}, x)$ would need a separate conversion. No procedure for learning the tensor from noisy data is asserted here.

IV.67 Source 63's Section 6: the face cover at Proposition 17.7 [write]

Added in the eleventh write, 7 October 2026.

The next section is source 63's Section 6, on the localized phase removal of Gowers's Proposition 17.7. *Logically it continues source 16's chapter* (Section IV.10), whose sliding windows and local ceiling it uses; it is printed here because no number printed earlier may move. Its conventions are in source 63's Section 2 (Section III.95), its introduction is Section F.68.

What is proved. Under the hypotheses of the local problem, for an odd box width m and prime $N \geq m^2$, a phase of degree at most $k + 1$ and a partition into progressions of lengths $\lceil m/(3k) \rceil$ or one more retain, as $m \rightarrow \infty$, the coefficient v_k/D_k with $v_k = 2^{k+1}/(k+2)!$ and $1 \leq D_k \leq e^{2/(3(k+2))}$ (Theorem IV.68.1), with explicit finite rounding factors (Theorem IV.68.5, Corollary IV.68.6); no coefficient in the class exceeds v_k (Proposition IV.68.7), so the best local coefficient is asymptotic to v_k . The new ingredient is an exact convolution identity for positive measures on the faces of a box (Theorem IV.68.2), extended to every positive real power of the kernel (Theorem IV.68.8).

Against source 16, and the predecessor that was never delivered (checked at the write). Source 16's window theorem and its local ceiling $2^{k+1}/(k+2)!$ (Proposition IV.13.5) are the sliding-window method and the ceiling used here; source 63 does not cite source 16. It credits instead the "preceding sliding-window localization coefficient" $v_k/\widehat{D}_k$, $\widehat{D}_k \rightarrow 1.2415870899$ (a retained fraction about 0.80542 of the ceiling), to a manuscript *Local Phase Removal and Cubic Rigidity in Gowers' Szemerédi Framework* of 7 October 2026, which was never delivered to this repository (Section F.4); that coefficient is not in this report and cannot be checked here. Against source 16's printed coefficients, which retain a fraction of v_k tending to zero (exponentially, at the optimized width $(2/(k+2))^{k+2}$), source 63's v_k/D_k closes the gap to first order; this answers source 16's question on the best local coefficient asymptotically (item IV.7), not at any fixed degree, as source 63 says.

Relation to the formal statements (checked at HEAD fb9602e55). `proposition_17_7` (Sections 17_18.lean, line 140) has the exact companion `proposition_17_7_holds` (Proofs17Localize dPhaseRemoval.lean, line 2941), with the paper's coefficient; nothing in this chapter is formalized.

IV.68 Asymptotically sharp local phase removal

Source 63, delivered Section 6.

The next result concerns the local phase-removal step of Gowers' Proposition 17.7 [1]. The preceding work [259] replaced the enumeration of carry patterns by sliding windows and a positive covering of the input increment box. It obtained a coefficient exceeding 80.54% of the geometric ceiling, uniformly in the degree. We improve the covering measure. The resulting coefficient is asymptotic to the ceiling as the degree tends to infinity. The new ingredient is an exact convolution identity for measures on the faces of a box; the phase polarization, sliding-window argument, and positive-cover interface are credited to the preceding work and are included here to make the proof self-contained.

IV.68.1 The local problem and its normalization

Let $G = \mathbb{Z}/N\mathbb{Z}$, where N is prime, and write $e_N(t) = \exp(2\pi it/N)$. If $u : G \rightarrow \mathbb{C}$ and $r \geq 1$, put

$$C_r(u) = \sum_{s \in G} \sum_{a \in G^r} \prod_{\omega \in \{0,1\}^r} \mathcal{C}^{|\omega|} u(s - \omega \cdot a) = N^{r+1} \|u\|_{U^r(G)}^{2^r}, \quad (\text{IV.68.1})$$

where $\mathcal{C}$ denotes complex conjugation. In particular these cube sums are nonnegative real numbers. For a partition $\mathcal{Q}$ into M progressions and an upper bound U for their lengths, define

$$B_k(u; \mathcal{Q}, U) = \frac{1}{MU^{k+2}} \sum_{Q \in \mathcal{Q}} C_{k+1}(\mathbf{1}_Q u). \quad (\text{IV.68.2})$$

This is the upper-cell-size normalization used in the corrected local interface. The factor U^{k+2} is retained throughout.

Here is the input. Let $m = 2h + 1$, $m^2 \leq N$, $d \in G^\times$, and

$$P_i = \{r_i + da : -h \leq a \leq h\}, \quad P = P_1 \times \cdots \times P_k.$$

Let $f : G \rightarrow \mathbb{C}$ satisfy $|f| \leq 1$. We explicitly assume that the coordinate function $a \mapsto \sigma(r + da)$ extends to a multi-affine polynomial on G^k , that is, a polynomial of degree at most one in each variable separately. The derivative-energy hypothesis is

$$\sum_{x \in P} \left| \sum_{s \in G} e_N(-s\sigma(x)) \prod_{\omega \in \{0,1\}^k} \mathcal{C}^{|\omega|} f(s - \omega \cdot x) \right|^2 \geq \alpha N^2 m^k, \quad \alpha > 0. \quad (\text{IV.68.3})$$

We seek a polynomial ϕ of degree at most $k + 1$ and a partition such that $B_k(fe_N(-\phi); \mathcal{Q}, U)$ is large.

Set

$$v_k = \frac{2^{k+1}}{(k+2)!}, \quad D_k = \sum_{j=0}^k \binom{k}{j} \left(\frac{2}{3k}\right)^j \frac{(k-j+2)!}{(k+2)!}. \quad (\text{IV.68.4})$$

Theorem IV.68.1 (The asymptotically sharp local coefficient). *Fix $k \geq 1$. As $m \rightarrow \infty$ through odd integers, uniformly over primes $N \geq m^2$ and inputs satisfying (IV.68.3), there are a polynomial ϕ of degree at most $k + 1$ and a partition into proper modular progressions with lengths*

$$L \quad \text{or} \quad L + 1, \quad L = \left\lceil \frac{m}{3k} \right\rceil,$$

such that

$$\mathbf{B}_k(fe_N(-\phi); \mathcal{Q}, L+1) \geq \alpha \left(\frac{v_k}{D_k} + O_k(m^{-1}) \right). \quad (\text{IV.68.5})$$

The $O_k(m^{-1})$ term describes the explicit finite coefficient in Theorem IV.68.5; it is uniform in the input functions and in $N \geq m^2$. Every cell has length between $m/(3k)$ and m . Moreover

$$1 + \frac{2}{3(k+2)} \leq D_k \leq \exp\left(\frac{2}{3(k+2)}\right), \quad D_k = 1 + \frac{2}{3k} + O(k^{-2}). \quad (\text{IV.68.6})$$

If c_k^{loc} is the best asymptotic coefficient when cells are required to have lengths between $m/(3k)$ and m , with $U = \max_{Q \in \mathcal{Q}} |Q|$, then

$$\frac{v_k}{D_k} \leq c_k^{\text{loc}} \leq v_k, \quad c_k^{\text{loc}} \sim \frac{2^{k+1}}{(k+2)!} \quad (k \rightarrow \infty). \quad (\text{IV.68.7})$$

The two limits are ordered: the local-scale limit $m \rightarrow \infty$ is taken first for each fixed k .

[write] **7 October 2026; sources 60–69.** Checked at the write (exact rationals). $D_1 = \widehat{D}_1 = 11/9$ and $D_2 = 127/108$, so $1/D_2 = 108/127 = 0.85039$; $D_k \leq \widehat{D}_k$ for every $k \leq 40$; $\widehat{D}_{40} = 1.24\dots$ Against source 16 (Section IV.10): source 16's optimized width retains $(2/(k+2))^{k+2}$, which is $2(k+2)!/(k+2)^{k+2}$ times v_k , a fraction decreasing to zero ($4/9$ at $k=1$, about $1.07 \cdot 10^{-4}$ at $k=10$); this theorem retains $v_k/D_k \geq e^{-2/(3(k+2))} v_k$. The ceiling v_k is source 16's Proposition IV.13.5, in the same normalization (every cell bounded by the normalizing size); source 16 is not cited here (Section IV.67).

In the definition of c_k^{loc} , “asymptotic coefficient c ” means that for every $\varepsilon > 0$ the conclusion with $\alpha(c - \varepsilon)$ holds for all sufficiently large m , uniformly over the remaining data. The theorem proves its first-order behavior as k grows. It does not determine the exact value at a fixed degree, or the relative $1/k$ correction to that optimum.

IV.68.2 An exact positive covering measure on box faces

For $w > 0$ put

$$K_w(x) = (w - \|x\|_1)_+^2, \quad x \in \mathbb{R}^k, \quad t_+ = \max(t, 0).$$

The elementary simplex integral gives

$$\int_{\mathbb{R}^k} K_w(x) dx = v_k w^{k+2}. \quad (\text{IV.68.8})$$

For $0 \leq j \leq k$, define

$$b_{k,j}(w) = \frac{(k-j+2)!}{(k+2)!} w^j. \quad (\text{IV.68.9})$$

On each codimension- j coordinate face of $B = [0, \ell]^k$, place constant density $b_{k,j}(w)$ relative to its $(k-j)$ -dimensional Lebesgue measure. Add all these measures, including the whole box for $j=0$, and denote the result by $\mu_{k,\ell,w}$. Face intersections receive the sum of their contributions. Its total mass is

$$T_k(\ell, w) = \sum_{j=0}^k \binom{k}{j} 2^j \ell^{k-j} \frac{(k-j+2)!}{(k+2)!} w^j. \quad (\text{IV.68.10})$$

Theorem IV.68.2 (Exact face-cover identity). *Assume $0 < w \leq \ell/2$. For $a \in B$, let $d_i = \min(a_i, \ell - a_i)$ and $d_S = \sum_{i \in S} d_i$. Then*

$$\int_B K_w(a - c) d\mu_{k,\ell,w}(c) = \frac{1}{(k+2)!} \sum_{S \subseteq [k]} 2^{k-|S|+1} d_S^{|S|} (w - d_S)_+^{k-|S|+2}. \quad (\text{IV.68.11})$$

The empty-set summand is interpreted as $2^{k+1}w^{k+2}$ before division by $(k+2)!$. In particular,

$$\int_B K_w(a - c) d\mu_{k,\ell,w}(c) \geq v_k w^{k+2} \quad (a \in B). \quad (\text{IV.68.12})$$

Proof. In each coordinate, only the nearest face can give a nonzero boundary contribution: the opposite face has distance at least $\ell/2 \geq w$. A tie at distance w contributes zero. Fix a set J of j nearest faces. The integral of the kernel over their intersection, before multiplication by the face density, is

$$I_J(a) = \frac{1}{(k-j+2)!} \sum_{S \subseteq [k] \setminus J} (-1)^{|S|} 2^{k-j-|S|+1} (w - d_{J \cup S})_+^{k-j+2}. \quad (\text{IV.68.13})$$

To verify this formula, begin with unrestricted integration in the $k-j$ free coordinates. The integral over a positive orthant is

$$\int_{[0,\infty)^r} (t - x_1 - \cdots - x_r)_+^2 dx = \frac{2t_+^{r+2}}{(r+2)!}.$$

Each unrestricted coordinate has two signs. Excluding a coordinate beyond its nearest boundary leaves one sign and shifts its distance variable by d_i . Inclusion–exclusion over the excluded coordinates therefore gives (IV.68.13).

Multiply by (IV.68.9) and sum over J . The factorial depending on j cancels. Group the terms by $T = J \cup S$, put $t = |T|$, and set $u = (w - d_T)_+$. If $u = 0$, every grouped term vanishes. If $u > 0$, the grouped contribution is

$$\begin{aligned} & \frac{2^{k-t+1}}{(k+2)!} \sum_{j=0}^t \binom{t}{j} (-1)^{t-j} w^j u^{k-j+2} \\ &= \frac{2^{k-t+1}}{(k+2)!} u^{k-t+2} (w - u)^t = \frac{2^{k-t+1}}{(k+2)!} (w - d_T)_+^{k-t+2} d_T^t. \end{aligned}$$

This proves (IV.68.11). The empty set supplies the full mass (IV.68.8); every other summand is nonnegative. At a corner all the other summands vanish, so the covering level is exact there. $\square$

The dependence of (IV.68.9) on the ambient dimension is decisive. For example the density on a facet is $w/(k+2)$. The previous exterior-point charging argument used $w/3$ at leading order. Integrating in the tangential coordinates before charging the boundary replaces that dimension-independent weight by the smaller one.

Lemma IV.68.3 (A lattice cover with no interpolation loss). *Let $h \geq 1$ be an integer, $A = [-h, h]^k \cap \mathbb{Z}^k$, and $0 < w \leq h$. There is a nonnegative measure ν on A of total mass $T_k(2h, w)$ such that*

$$\sum_{c \in A} \nu(c) K_w(a - c) \geq v_k w^{k+2} \quad (a \in A). \quad (\text{IV.68.14})$$

Explicitly, put $q(c) = 1$ for $-h < c < h$ and $q(\pm h) = 1/2$, and let $B(c) = \{i : c_i \in \{-h, h\}\}$. Then one can take

$$\nu(c) = \sum_{J \subseteq B(c)} b_{k,|J|}(w) \prod_{i \notin J} q(c_i). \quad (\text{IV.68.15})$$

Proof. Translate Theorem IV.68.2 to $[-h, h]^k$. Fix an integer target a . On any closed unit interval with integer endpoints in a center coordinate c_i , the function $|a_i - c_i|$ is affine. Thus $K_w(a - c)$ is convex in that coordinate on the interval, because $t \mapsto (t_+)^2$ is convex. It follows by successive interpolation in all coordinates that rounding a continuous center to adjacent lattice endpoints, with independent roundings preserving each coordinate's mean, can only increase the expected kernel value.

This single rounding probability kernel is independent of a . Apply it to the continuous face measure. It preserves total mass and increases the convolution simultaneously at every integer target. An interior integer receives integrated weight one from its two adjacent unit intervals; an endpoint receives one-half from its sole adjacent interval. A coordinate fixed on a face is unchanged. These observations give (IV.68.15), and (IV.68.14) follows from (IV.68.12). Summing the weights also gives (IV.68.10) with $\ell = 2h$ directly. $\square$

Global convexity of K_w is neither asserted nor needed. Convexity on each coordinate interval with integer endpoints is the property that makes the simultaneous rounding valid.

IV.68.3 From any positive cover to local cube mass

We give the precise analytic interface used to insert the new cover. The mechanism in this subsection is the preceding report's positive-cover argument [259].

Let $N > k + 1$ be prime, and let $\tau : G^k \rightarrow G$ be multi-affine. For functions $f_\epsilon : G \rightarrow \mathbb{C}$ indexed by $\epsilon \in \{0, 1\}^k$, write

$$F(a) = \sum_{s \in G} e_N(-s\tau(a)) \prod_{\epsilon \in \{0, 1\}^k} C^{|\epsilon|} f_\epsilon(s - \epsilon \cdot a), \quad E_A = \sum_{a \in A} |F(a)|^2, \quad (\text{IV.68.16})$$

where $A \subset \mathbb{Z}^k$ is finite and integer arguments are reduced modulo N . No bound on the f_ϵ is required here.

Lemma IV.68.4 (Positive-cover localization). *Let $3 \leq M < N$, $w = N/M$, and $U = \lceil w \rceil$. Suppose a nonzero nonnegative measure ν on a finite set $C \subset \mathbb{Z}^k$ has total mass T and satisfies*

$$\sum_{c \in C} \nu(c) K_w(a - c) \geq \kappa \quad (a \in A).$$

Then for some $c \in C$, some $\epsilon \in \{0, 1\}^k$, a polynomial ϕ_ϵ^c of degree at most $k + 1$, and a partition $\mathcal{Q}$ into intervals of lengths $\lfloor w \rfloor$ or $\lceil w \rceil$, one has

$$\mathbf{B}_k(f_\epsilon(\cdot - \epsilon \cdot c) e_N(-\phi_\epsilon^c); \mathcal{Q}, U) \geq \frac{E_A \kappa}{N^2 T U^{k+2}}. \quad (\text{IV.68.17})$$

The polynomials are explicitly determined by the recentered multi-affine polynomial $\tau_c(b) = \tau(b + c)$.

Proof. Put $E_c = \sum_{a \in A} K_w(a - c)|F(a)|^2$. The cover implies $\sum_c \nu(c)E_c \geq \kappa E_A$, so choose c with $E_c \geq \kappa E_A/T$. Let

$$f_\epsilon^c(x) = f_\epsilon(x - \epsilon \cdot c), \quad \tau_c(b) = \tau(b + c).$$

The corresponding expression satisfies $F^c(b) = F(b + c)$.

Embed G into $\mathbb{Z}/(NM)\mathbb{Z}$ by $x \mapsto Mx$ and define windows

$$S_t = \{x \in G : (Mx - t) \bmod NM \in \{1, \dots, N\}\}, \quad t \in \mathbb{Z}/(NM)\mathbb{Z}.$$

Every S_t is a cyclic interval of size $\lfloor w \rfloor$ or $\lceil w \rceil$. For $0 \leq j < N$, the family

$$\mathcal{P}_j = (S_{j+\ell N})_{0 \leq \ell < M} \tag{IV.68.18}$$

is a partition: the corresponding M blocks of N consecutive residues partition the larger cyclic group.

If $b \in \mathbb{Z}^k$ has $D(b) = \|b\|_1 < w$, the integer span of the cube vertices $s - \epsilon \cdot b$ is exactly $D(b)$. Exactly $N - MD(b)$ windows contain all the vertices. In the embedded circle their span is $MD(b) < N$; the possible starts of a containing block form an interval with $N - MD(b)$ integers. There is no additional containing arc, because $D(b) < w < N/2$ makes the complementary gap longer than the block length.

Let $F_t^c(b)$ denote $F^c(b)$ with each vertex function restricted to S_t . The window count gives

$$\sum_t F_t^c(b) = (N - MD(b))F^c(b).$$

The set $\mathcal{D} = (A - c) \cap \{b : \|b\|_1 < w\}$ embeds injectively into G^k , since $2w < N$. Cauchy-Schwarz over the NM windows therefore gives

$$\begin{aligned} M^2 E_c &\leq NM \sum_t \sum_{b \in \mathcal{D}} |F_t^c(b)|^2 \\ &\leq NM \sum_t \sum_{b \in G^k} |F_t^c(b)|^2. \end{aligned} \tag{IV.68.19}$$

The domain is enlarged while every summand is a nonnegative square.

For completeness, write $\tau_c(b) = \sum_{J \subseteq [k]} d_J \prod_{i \in J} b_i$ and set

$$\phi_\epsilon^c(X) = \sum_{J \supseteq \text{supp}(\epsilon)} \frac{d_J}{(|J| + 1)!} X^{|J|+1}. \tag{IV.68.20}$$

All factorials are invertible modulo N . The finite-difference identity

$$\sum_{\omega \in \{0,1\}^r} (-1)^{|\omega|} (s - \omega \cdot z)^r = r! \prod_{i=1}^r z_i$$

applied separately to each J yields

$$y \tau_c(b) = \sum_{\epsilon \in \{0,1\}^k} \sum_{z=0}^1 (-1)^{|\epsilon|+z} \phi_\epsilon^c(s - \epsilon \cdot b - zy). \tag{IV.68.21}$$

Put $g_\epsilon^c = f_\epsilon^c e_N(-\phi_\epsilon^c)$. Expand a square in (IV.68.19), writing its second base variable as $s - y$. Its phase is $e_N(-y\tau_\epsilon(b))$. Identity (IV.68.21) absorbs this phase into the vertex functions. The mixed Gowers–Cauchy–Schwarz inequality gives

$$\sum_{b \in G^k} |F_t^c(b)|^2 \leq \prod_{\epsilon \in \{0,1\}^k} C_{k+1}(\mathbf{1}_{S_t} g_\epsilon^c)^{1/2^k}. \quad (\text{IV.68.22})$$

Each first- k -coordinate vertex occurs twice with the same function, which explains the exponent 2^{-k} .

Hölder’s inequality in t and (IV.68.19) imply

$$\prod_{\epsilon} \left(\sum_t C_{k+1}(\mathbf{1}_{S_t} g_\epsilon^c) \right)^{1/2^k} \geq \frac{M}{N} E_c.$$

One factor is at least $(M/N)E_c$. Fix that ϵ . Every window occurs exactly once among the N partitions (IV.68.18), so some partition has cube sum at least $(M/N^2)E_c$. Divide by MU^{k+2} and use $E_c \geq \kappa E_A/T$ to obtain (IV.68.17). $\square$

IV.68.4 The finite theorem and its sharp asymptotic consequence

Theorem IV.68.5 (Explicit finite coefficient). *Assume (IV.68.3), $m = 2h + 1 > 3k$, and $m^2 \leq N$. Put*

$$L = \left\lceil \frac{m}{3k} \right\rceil, \quad M = \left\lfloor \frac{N}{L} \right\rfloor, \quad w = \frac{N}{M}, \quad (\text{IV.68.23})$$

and assume $w \leq h$. There are a polynomial ϕ of degree at most $k+1$ and a partition into progressions of lengths L or $L+1$, with the same common difference d as the input, such that

$$B_k(f e_N(-\phi); \mathcal{Q}, L+1) \geq \alpha R_{k,m,N}, \quad R_{k,m,N} = v_k \frac{m^k w^{k+2}}{T_k(m-1, w)(L+1)^{k+2}}. \quad (\text{IV.68.24})$$

The additional condition $w \leq h$ is automatic when $k \geq 2$, and when $k = 1, m \geq 9$.

Proof. Put $\tilde{f}(t) = f(dt)$, $q_i = d^{-1}r_i$, and

$$f_\epsilon(t) = \tilde{f}(t - \epsilon \cdot q), \quad \tau(a) = d\sigma(r + da).$$

The polynomial τ is multi-affine. Substituting $s = dt$ into each Fourier sum in (IV.68.3) gives exactly (IV.68.16), with these functions. Thus their input energy is at least $\alpha N^2 m^k$ on $A = [-h, h]^k \cap \mathbb{Z}^k$.

One has $2 \leq L \leq m$, $M \geq L$, and $3 \leq M < N$. Writing $N = LM + r$, $0 \leq r < L$, gives $L \leq w < L+1$. In fact $r > 0$, since N is prime and $2 \leq L < N$, so $\lfloor w \rfloor = L$ and $\lceil w \rceil = L+1$. Apply Lemma IV.68.3 and then Lemma IV.68.4, with

$$T = T_k(m-1, w), \quad \kappa = v_k w^{k+2}.$$

This gives (IV.68.24) for one of the functions $\tilde{f}(s - \epsilon \cdot (q + c)) e_N(-\phi_\epsilon^c(s))$. If $b = \epsilon \cdot (q + c)$, translate the variable by b . The resulting function is $\tilde{f}(s) e_N(-\phi_\epsilon^c(s + b))$, and the support partition is translated by $-b$. Polynomial degree is preserved. Reversing the initial dilation gives the original function $f(x)$ twisted by $\phi(x) = \phi_\epsilon^c(d^{-1}x + b)$, on progressions of common difference d . This again preserves the polynomial degree.

For the claimed width condition, $L + 1 \leq h$ is equivalent to $(3k - 2)h \geq 3k + 1$. It holds when $k = 1, h \geq 4$, and for every $k \geq 2$ in the range $m > 3k$. Since $w < L + 1$, this is sufficient. Finally, $L \geq m/(3k)$ and $L + 1 \leq m/3 + 2 \leq m$ in the stated range, so every output length has the required bounds. $\square$

For the two excluded pairs $(k, m) = (1, 5), (1, 7)$, or for the branch $m \leq 3k$, the preceding finite localization results remain available. Those finitely many scales do not affect Theorem IV.68.1. The new finite expression is not claimed to dominate every earlier finite expression at the smallest parameters.

Corollary IV.68.6 (Uniform finite bounds and joint scale limits). *For the finite coefficient in Theorem IV.68.5,*

$$\exp\left(-\frac{k+2}{L} - \frac{2kw}{(k+2)(m-1)}\right) \leq \frac{R_{k,m,N}}{v_k} \leq \exp\left(\frac{k}{m-1}\right). \quad (\text{IV.68.25})$$

More precisely, let $k \rightarrow \infty$ along any sequence of admissible parameters. If $m/k^2 \rightarrow \lambda \in [0, \infty]$, then, uniformly over the choices of prime $N \geq m^2$,

$$\frac{R_{k,m,N}}{v_k} \longrightarrow \begin{cases} 0, & \lambda = 0, \\ e^{-3/\lambda}, & 0 < \lambda < \infty, \\ 1, & \lambda = \infty. \end{cases} \quad (\text{IV.68.26})$$

In particular, this finite construction has relative retention tending to one whenever $m/k^2 \rightarrow \infty$. The assertion concerns the explicit coefficient $R_{k,m,N}$, and does not assert that this scale condition is necessary for the unknown optimal finite theorem.

Proof. Write $\ell = m - 1$ and

$$H_k(z) = \sum_{j=0}^k \binom{k}{j} \frac{(k-j+2)!}{(k+2)!} z^j.$$

Equation (IV.68.29), whose elementary product proof does not depend on the choice of z , gives

$$1 \leq H_k(z) \leq e^{kz/(k+2)} \quad (z \geq 0).$$

The exact expression for the coefficient is

$$\frac{R_{k,m,N}}{v_k} = \left(\frac{m}{\ell}\right)^k \left(\frac{w}{L+1}\right)^{k+2} H_k(2w/\ell)^{-1}. \quad (\text{IV.68.27})$$

Use $1 \leq (m/\ell)^k \leq e^{k/\ell}$, $L/(L+1) \leq w/(L+1) < 1$, and $(L/(L+1))^{k+2} \geq e^{-(k+2)/L}$ to obtain (IV.68.25).

The partition rounding is particularly small. Since $m > 3k$, $L < 2m/(3k)$; also $M \geq N/(2L)$. Hence, uniformly over $N \geq m^2$,

$$0 < w - L < \frac{L}{M} \leq \frac{2L^2}{N} < \frac{8}{9k^2}. \quad (\text{IV.68.28})$$

If $0 < \lambda \leq \infty$, then $m/k \rightarrow \infty$ and $L \sim m/(3k)$. In the logarithm of (IV.68.27), the first factor contributes $k \log(m/\ell) = o(1)$ and the last contributes $0 \leq \log H_k(2w/\ell) \leq 2kw/((k+2)\ell) = o(1)$.

Furthermore (IV.68.28) gives

$$\begin{aligned} (k+2) \log \frac{w}{L+1} &= -(k+2) \log(1+L^{-1}) + (k+2) \log \left(1 + \frac{w-L}{L}\right) \\ &= -\frac{k+2}{L} + O(k/L^2) + O(1/(kL)) \rightarrow -\frac{3}{\lambda}, \end{aligned}$$

with $3/\infty = 0$. This proves the positive and infinite cases.

If $\lambda = 0$, then $L = o(k)$, while $k/(m-1) \leq 1/3$. For sufficiently large k , (IV.68.28) gives $w \leq L + 1/2$. Therefore

$$0 \leq \frac{R_{k,m,N}}{v_k} \leq e^{1/3} \left(1 - \frac{1}{2(L+1)}\right)^{k+2} \leq \exp\left(\frac{1}{3} - \frac{k+2}{2(L+1)}\right) \rightarrow 0.$$

All estimates are uniform in the ambient prime, completing the proof. $\square$

Proof of the lower bound in Theorem IV.68.1. For fixed k , uniformly for $N \geq m^2$, one has

$$\frac{w}{m} = \frac{1}{3k} + O_k(m^{-1}), \quad \frac{w}{L+1} = 1 + O_k(m^{-1}), \quad \frac{m-1}{m} = 1 + O(m^{-1}).$$

Substitution into (IV.68.10) gives

$$\frac{T_k(m-1, w)}{m^k} = D_k + O_k(m^{-1}).$$

Equation (IV.68.24) therefore proves (IV.68.5). The first two terms of D_k give the lower bound in (IV.68.6). For the upper bound, observe that

$$\binom{k}{j} \frac{(k-j+2)!}{(k+2)!} = \frac{1}{j!} \prod_{r=0}^{j-1} \frac{k-r}{k+2-r} \leq \frac{1}{j!} \left(\frac{k}{k+2}\right)^j. \quad (\text{IV.68.29})$$

Comparison with the exponential series proves $D_k \leq e^{2/[3(k+2)]}$ and the stated expansion. $\square$

Proposition IV.68.7 (The local geometric ceiling). *In the class defining c_k^{loc} , no asymptotic coefficient exceeds v_k .*

Proof. Use $f \equiv 1$, $\sigma \equiv 0$, and $\alpha = 1$ in (IV.68.3). Every polynomial twist has modulus one, so the triangle inequality in the cube expansion gives

$$0 \leq C_{k+1}(\mathbf{1}_Q f e_N(-\phi)) \leq C_{k+1}(\mathbf{1}_Q).$$

Let $A_r(t)$ be the number of integer r -cubes contained in an interval of length t . Their span is the sum of the absolute increments, whence

$$A_r(t) = \sum_{a \in \mathbb{Z}^r} (t - \|a\|_1)_+, \quad \sum_{t \geq 0} A_r(t) z^t = \frac{z(1+z)^r}{(1-z)^{r+2}}. \quad (\text{IV.68.30})$$

Coefficient extraction yields

$$A_r(t) = \frac{2^r}{(r+1)!} t^{r+1} + O_r(t^r).$$

For a proper progression $Q \subset G$, an invertible dilation identifies it with an interval. If $N > 2r|Q|$, every cube with all vertices in that interval has consistent integer lifts: the congruences between its lifted vertices differ by less than N , and hence are integer equalities. Thus its modular cube count is $A_r(|Q|)$. For fixed k , this condition holds eventually uniformly when $|Q| \leq m$ and $N \geq m^2$.

Take $r = k + 1$. All cell sizes satisfy $|Q| \geq m/(3k)$, so

$$\mathbf{C}_{k+1}(\mathbf{1}_Q) \leq (v_k + O_k(m^{-1}))|Q|^{k+2}.$$

Sum over cells and divide by MU^{k+2} , using $|Q| \leq U$. The quotient is at most $v_k + O_k(m^{-1})$, proving the ceiling. $\square$

The proposition and the lower bound prove (IV.68.7). In particular, the theorem closes the constant-factor asymptotic gap left by the previous positive cover. It establishes neither an exact fixed-degree optimum nor a new global bound for Szemerédi's theorem.

IV.68.5 Comparison with the earlier covering coefficient

The previous coefficient [259] was $v_k/\widehat{D}_k$, where

$$\widehat{D}_k = \sum_{j=0}^k \binom{k}{j} \left(\frac{2}{3k}\right)^j \frac{2}{(j+2)!}, \quad \widehat{D}_k \longrightarrow 1.2415870899 \dots$$

For $0 \leq j \leq k$,

$$\frac{(k-j+2)!}{(k+2)!} \leq \frac{2}{(j+2)!},$$

because $\binom{k+2}{j} \geq \binom{j+2}{j}$. It follows that $D_k \leq \widehat{D}_k$, with strict inequality for $k \geq 2$. At $k = 1$ both denominators are $11/9$. Some retained fractions of the common ceiling are listed below; the checks in the accompanying script use rational arithmetic before decimal presentation.

k	previous $1/\widehat{D}_k$	new $1/D_k$	exact new value where short
1	0.81818	0.81818	9/11
2	0.81203	0.85039	108/127
3	0.80988	0.87648	10935/12476
4	0.80879	0.89540	93312/104213
8	0.80712	0.93556	
20	0.80610	0.97015	
100	0.80556	0.99349	

The denominator has the useful equivalent expression

$$D_k = \frac{1}{(k+1)(k+2)} \sum_{j=0}^k \frac{(k-j+1)(k-j+2)}{j!} \left(\frac{2}{3k}\right)^j. \quad (\text{IV.68.31})$$

For example, expanding its first four terms and bounding the positive remainder gives

$$D_k = 1 + \frac{2}{3k} - \frac{10}{9k^2} + \frac{148}{81k^3} + O(k^{-4}),$$

and hence

$$\frac{1}{D_k} = 1 - \frac{2}{3k} + \frac{14}{9k^2} - \frac{292}{81k^3} + O(k^{-4}).$$

These are expansions of the proved lower coefficient, not of the unknown optimum c_k^{loc}/v_k .

IV.68.6 Other powers of the kernel

The positive identity extends beyond the square kernel. This geometric generalization is independent of its present application to Proposition 17.7.

Theorem IV.68.8 (Positive-power face covers). *Let $s > 0$, $K_{w,s}(x) = (w - \|x\|_1)_+^s$, and $0 < w \leq \ell/2$. On every codimension- j face of $[0, \ell]^k$ place density*

$$b_{k,j,s}(w) = \frac{\Gamma(k - j + s + 1)}{\Gamma(k + s + 1)} w^j,$$

and let μ_s be the sum of these measures. With d_S as in Theorem IV.68.2,

$$\int K_{w,s}(a - c) d\mu_s(c) = \frac{\Gamma(s + 1)}{\Gamma(k + s + 1)} \sum_{S \subseteq [k]} 2^{k-|S|} d_S^{|S|} (w - d_S)_+^{k-|S|+s}. \quad (\text{IV.68.32})$$


In particular this is at least $2^k \Gamma(s + 1) w^{k+s} / \Gamma(k + s + 1)$, the full integral of the kernel. If $s \geq 1$, the same cover passes to integer centers by the interpolation procedure of Lemma IV.68.3, with unchanged mass and the corresponding face weights.

Proof. Replace the simplex integral in (IV.68.13) by

$$\int_{[0,\infty)^r} (t - x_1 - \cdots - x_r)_+^s dx = \frac{\Gamma(s + 1)}{\Gamma(r + s + 1)} t_+^{r+s}.$$

The face density cancels the denominator exactly. The same inclusion–exclusion grouping leaves $u^{k-|S|+s} (w - u)^{|S|}$, proving (IV.68.32). All exponents at $u = 0$ are positive because $s > 0$. For the lattice assertion, $t \mapsto (t_+)^s$ is convex for $s \geq 1$, so the separate-cell convexity and mean-preserving rounding proof is unchanged. $\square$

The exact verification file checks the lattice cover and its mass at 10,980 integer targets in 32 parameter cases. These checks exercise boundary intersections, half-integer widths, and the endpoint weights. They support implementation of the formulas; the general results follow from the proofs above.



Proved lower coefficients; the limit $m \rightarrow \infty$ is taken first for each k .

Figure IV.68.1: Fractions of the geometric ceiling retained by the previous and new localization coefficients. The new fraction $1/D_k$ tends to one; the preceding fraction tends to approximately 0.80542. The horizontal line is the universal geometric upper bound.

Part V

Lean interface, ledger crosswalk and research questions

Contents. Sections V.1 and V.2 are the write’s: the crosswalk of every comparison with the formal project, checked against the current ledger, and the deduplicated list of the sources’ 46 research questions with credits (Vladimir’s standing rule of 4 October 2026). Then each source’s remaining sections, in source order: 01’s formalization interface, questions, conclusions and appendices (delivered Sections 14–16, Appendices A–B); 02’s (Sections 10–12, Appendix A); 03’s iteration bookkeeping, Lean route, questions and reproducibility (Sections 9–11, Appendix A); 04’s quantitative comparisons, questions, closing perspective and appendices (Sections 9–11, Appendices A–B); 05’s Lean relation, computation, questions and conclusion (Sections 7–10). The sources’ question sections are printed as delivered, with dated notes where another source answers a question.

Reading conventions of this Part (the full dictionary is Section F.3).

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
pins, counts	01, 03: pin 3a25020a6; 04, 05: pin 327773149; 02: observed branch reference 66c02f8a1. The ledger counts quoted by the sources (88 exact, 32 open) are those of the pins; at HEAD they are 98 and 22 (Section V.1).
Qk , (Q1), (Q2)	04: its questions are numbered Q1–Q12 and two displays are tagged (Q1), (Q2); 01: numbered Questions (Section V.4); 03: numbered Research questions (Section V.14); 02, 05: subsections and paragraphs. Item names in Section V.2 are the write’s.

Added in the second write (6 October 2026): sources 06–10. Contents. The write’s Sections V.1 and V.2 are extended by dated additions: the crosswalk to the ledger at HEAD of this write, and the 50 research questions of sources 06–10 (plus two inline question paragraphs of source 10) merged into the theme lists. Then the remaining sections of the new sources, in source order: 06’s iteration potential, Lean map and questions (delivered Sections 6, 8, 9); 07’s verification, questions, conclusion and reproduction commands (Sections 9–11, Appendix B); 08’s Lean integration, checks, questions, conclusion and build provenance (Sections 11–14, Appendix B); 09’s checks, Lean integration, questions, conclusion and dependency map (Sections 11, 12, 14, 15, Appendix C); 10’s iteration bookkeeping, Lean map, questions and provenance (Sections 10–12, Appendix A).

Reading conventions for the sections of sources 06–10 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
pins, counts (06–10)	06, 09, 10: pin 7d99e5dee; 07: pin ae0e89574; 08: no commit, blob identifiers of 7d99e5dee. The ledger counts quoted by 06 and 10 (95 exact, 25 open) and 06’s “140 modules, 1,365 theorems” are those of 7d99e5dee; at HEAD 764740f07 they are 99 and 21, and FORMALIZATION_STATUS.txt reports 196 modules and 1,805 audited theorems (Section V.1).
iteration	06: F , G the increment and the charge of Lemma V.25.1, c, p, q, a_0 the local rule, $A = \log(1/a_0)$, $L = \log(1/\delta_0)$, $B(\delta_0)$ the double-logarithmic budget. 10: c , a the increment rule, ϑ_i , b_i the length rule, P_t , M_t . 06’s L is a <i>logarithm</i> , not a <i>length</i> ; p is an <i>exponent</i> , not a <i>prime</i> .

Added in the third write (6 October 2026): sources 11–14. Contents. The write’s Sections V.1 and V.2 gain dated additions: the crosswalk of sources 11–14 at HEAD of this write and their 45 questions

merged into the theme lists. Then the remaining sections of the new sources, in source order: 11’s verification and formalization interfaces, questions, conclusion and reproduction (delivered Sections 12–14, Appendix A); 12’s formalization interfaces and verification, questions and parameter ledger (Sections 10–11, Appendix B); 13’s finite examples, verification and formalization, questions and reproduction (Sections 10–11, Appendix A); 14’s Lean source map, research program and reproducibility (Sections 8–9, Appendix B).

Reading conventions for the sections of sources 11–14 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
pins, counts (11–14)	11: pin 40f885098; 12, 13, 14: pin 9ea383e9b. The counts 12 and 13 quote (97 exact, 23 open; 154 modules, 1,508 theorems) are those of the pins; at HEAD e1db00e99 they are 100 and 20, 206 modules and 1,879 theorems (Section V.1).
questions	11 and 13: “Research question” <i>k.j</i> (here with Part numbers); 12: “Question”; 14: the numbered items 1–12 of its research program (“14 Q <i>k</i> ” in Section V.2).

Added in the fourth write (6 October 2026): sources 15–20. Contents. The write’s Sections V.1 and V.2 gain dated additions: the crosswalk of sources 15–20 at HEAD of this write, a re-check of the earlier tables, and their 74 questions merged into the theme lists. Then the remaining sections of the new sources, in source order: 15’s computations and Lean plan, questions, conclusion and appendices (delivered Sections 13–15, Appendices A–B); 16’s connections, source fidelity and plan, computations and questions (Sections 9–12, Appendix A); 17’s Lean relation, audits, questions, conclusion and appendices (Sections 9–12, Appendices A–B); 18’s quantitative consequences, Lean relation, verification, questions and source record (Sections 6–9, Appendix A); 19’s source audit, verification and questions (Sections 7, 9, 10); 20’s formalization, questions, conclusion and source audit (Sections 9–11, Appendix A).

Reading conventions for the sections of sources 15–20 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
pins, counts (15–20)	15: pin 9ea383e9b; 16–19: pin dca9f0638; 20: observed checkpoint 1bd57e91c. The counts 16 and 19 quote (98 exact, 22 open; 169 modules, 1,658 theorems) are those of dca9f0638; at HEAD 0b4890a1c they are 101 and 19, 233 modules and 2,138 theorems (Section V.1).
questions	15: Questions <i>k.1–k.9</i> ; 16, 18: numbered items; 17: numbered paragraphs; 19: Research questions and five items at the end of its Section 6 (“S6(1)”–“S6(5)”); 20: Questions numbered within their section here. In Section V.2 they are “15 Q1” and so on.

Added in the fifth write (6 October 2026): sources 21–27. Contents. The write’s Sections V.1 and V.2 gain dated additions: the crosswalk of sources 21–27 at HEAD of this write, and their 74 questions merged into the theme lists. Then the remaining sections of the new sources, in source order: 21’s repository interface, questions and appendices (delivered Sections 10–11, Appendices A–B); 22’s verification and questions (Sections 11–12); 23’s verification record, questions and build appendix (Sections 9–10, Appendix B); 24’s examples and checks, Lean integration, questions, conclusion and appendices (Sections 10–13, Appendices A–B); 25’s integration record, questions and conclusion (Sections 9–11); 26’s source comparison with its questions, and its companion (Sections 5–6); 27’s verification and integration, questions, conclusion and appendix (Sections 12–14, Appendix A).

Reading conventions for the sections of sources 21–27 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
pins (21–27)	21, 22: pin 43306653a; 23: 65f4e039f; 24: 2e0553b0c; 25: 568d1802e; 26: c94998a5d; 27: blob identifiers only. None of them quotes ledger counts; every Lean file they cite is unchanged at HEAD 2dfe46f44 (Section V.1).
questions	21, 24, 25: Questions; 22, 27: Research questions; 23, 26: numbered items. In Section V.2 they are “21 Q1” and so on, in delivered order.

Added in the sixth write (6 October 2026): sources 28–33. Contents. The write’s Sections V.1 and V.2 gain dated additions: the crosswalk of sources 28–33 at HEAD of this write, and their 69 questions merged into the theme lists. Then the remaining sections of the new sources, in source order: 28’s interface, computations, questions and provenance (delivered Sections 9–12); 29’s integration, verification and questions (Sections 10–12); 30’s crosswalk, questions, conclusion and build appendix (Sections 10–12, Appendix A); 31’s verification, questions and dictionary (Sections 10–11, Appendix B); 32’s source comparison and questions (Sections 6–7); 33’s questions, integration and audit checklist (Sections 9–10, Appendix A).

Reading conventions for the sections of sources 28–33 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
pins (28–33)	28, 29, 31: pin e3125e698; 30: 8d65d3c9f; 32: c94998a5d; 33: 9f83dbe0f. None of them quotes ledger counts; every Lean file they cite is unchanged at HEAD 5d4718772 (Section V.1).
questions	29, 30: Questions; 28, 31, 33: Research questions; 32: numbered items. In Section V.2 they are “28 Q1” and so on, in delivered order.

Added in the seventh write (6 October 2026): sources 34–39. Contents. The write’s Sections V.1 and V.2 gain dated additions: the crosswalk of sources 34–39 at HEAD of this write, and their 73 questions merged into the theme lists. Then the remaining sections of the new sources, in source order: 34’s consequences and research program, and integration (delivered Sections 11–12); 35’s evidence and formalization, and research programme (Sections 8–9); 36’s checks, integration, questions and conclusion (Sections 11–14); 37’s verification, questions, conclusion and source audit (Sections 10–12, Appendix C); 38’s verification, integration and questions (Sections 8–10); 39’s interface, checks, questions, conclusion, dependency map and notation checklist (Sections 12–15, Appendices A–B).

Reading conventions for the sections of sources 34–39 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
pins (34–39)	34, 35, 38: pin 12da35673; 36: 6c2e2172b; 37, 39: 869289aa3. Only source 38 quotes ledger counts (those of its pin); every Lean and paper file they cite is unchanged at HEAD 89c1cbc38 (Section V.1).
questions	34, 35, 36, 39: Research questions; 37: Questions; 38: numbered items. In Section V.2 they are “34 Q1” and so on, in delivered order (source 35’s question in its Section 6 is “35 Q1”).

Added in the eighth write (7 October 2026): sources 40–46. Contents. The write’s Sections V.1 and V.2 gain dated additions: the crosswalk of sources 40–46 at HEAD of this write, including the repository’s new formal work on Corollary 5.8, and their 62 questions merged into the theme lists. Then the remaining sections of the new sources, in source order: 40’s reproducibility, questions and conclusion (delivered Sections 10–12); 41’s pinned source and context, and reproducibility and questions (Sections 8–9); 42’s relation to prior arguments and questions, and diagnostics (Sections 8–9); 43’s relation to Gowers’s argument and the formal project, checks, questions and conclusion (Sections 10–13); 44’s verification and formalization

targets, questions, conclusion and reproduction appendix (Sections 11–13, Appendix A); 45’s research programme and its appendices on normalization, certificates and sources (Section 6, Appendices A–C); 46’s verification, questions, conclusion and dependency map (Sections 10–12, Appendix A).

Reading conventions for the sections of sources 40–46 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
pins (40–46)	40, 46: no pin (46 records two blobs); 41: <code>c94998a5d</code> ; 42: <code>b16ce380d</code> ; 43: <code>797aa3cfe</code> ; 44: <code>8c12c906e</code> ; 45: <code>6ce33539e</code> . Every Lean and paper file they cite is unchanged at HEAD <code>6d4942222</code> (Section V.1).
questions	40, 43, 46: Research questions; 44, 45: Questions; 41, 42: enumerated items. In Section V.2 they are “40 Q1” and so on, in delivered order.

Added in the ninth write (7 October 2026): sources 47–53. Contents. The write’s Sections V.1 and V.2 gain dated additions: the crosswalk of sources 47–53 at HEAD of this write, including the repository’s formal work on Lemma 16.10 and Section 13, and their 57 questions merged into the theme lists. Then the remaining sections of the new sources, in source order: 47’s relation to Gowers’s argument and the formal project, verification, questions, conclusion and reproduction appendix (delivered Sections 10–13, Appendix A); 48’s checks, integration, questions, conclusion and appendices on algorithms and dependencies (Sections 10–13, Appendices A–B); 49’s consequences, questions and reproducibility (Sections 9–11); 50’s checks, formalization route, questions, conclusion and appendices (Sections 10–13, Appendices A–C); 51’s interfaces after complete extraction, verification roles and source map (Sections 11–12, Appendix A); 52’s repairs and open questions, reproducibility, source map and reconstruction guide (Sections 10–11, Appendices A–B); 53’s scope and questions, companion and source map (Sections 9–10, Appendix A).

Reading conventions for the sections of sources 47–53 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
pins (47–53)	47, 52: <code>5c9a442d2</code> ; 48: <code>0b4890a1c</code> ; 49: <code>8c12c906e</code> ; 50: <code>588552b79</code> ; 51: <code>aad4789fc</code> , ..., <code>4f8cd1198</code> ; 53: <code>bf2fe146d</code> . Changes at HEAD <code>06b30ca52</code> to the files they cite are recorded in Section V.1.
questions	47, 48, 50: Research questions; 49: Questions (and one paragraph in its Section 8); 51: enumerated interface items; 52: paragraphs; 53: enumerated items. In Section V.2 they are “47 Q1” and so on, in delivered order.

Added in the tenth write (7 October 2026): sources 54–59. Contents. The write’s Sections V.1 and V.2 gain dated additions: the crosswalk of sources 54–59 at HEAD of this write, including the formal proof of Theorem 13.12, and their 63 questions merged into the theme lists. Then the remaining sections of the new sources, in source order: 54’s repository integration, questions and conclusion (delivered Sections 10–12); 55’s questions, validation and appendices on the coefficient ledger and formal Taylor expansions (Sections 12–13, Appendices A–B); 56’s decoder and diagnostics, Gowers interface, questions and proof audit (Sections 10–12, Appendix A); 57’s questions and proof dependencies (Sections 7–8); 58’s examples and checks, questions, integer certificate and package (Sections 12–13, Appendices A–B); 59’s verification scope and questions (Sections 11–12).

Reading conventions for the sections of sources 54–59 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
pins (54–59)	54: <code>6a6d7961c</code> ; 55, 57: <code>128f514af</code> ; 56: <code>b6a8ee966</code> ; 58: <code>5d4718772</code> ; 59: <code>bec460a33</code> . Changes at HEAD <code>fe7165a3c</code> to the files they cite are recorded in Section V.1.
questions	54, 55, 56: Questions 11.1–11.11, 12.1–12.10, 12.1–12.10; 57: Research questions 1–12; 58, 59: Questions 13.1–13.10, 12.1–12.10. In Section V.2 they are “57 Q1” and so on, in delivered order.

Added in the eleventh write (7 October 2026): sources 60–69. Contents. The write’s Sections V.1 and V.2 gain dated additions: the crosswalk of sources 60–69 at HEAD of this write, and their 104 questions merged into the theme lists. Then the remaining sections of the new sources, in source order: 60’s consequences, evidence, questions and decoder (Sections 10–12, Appendix A); 61’s checks, integration, questions, conclusion and ledger (Sections 10–13, Appendix A); 62’s checks, integration, questions and conclusion (Sections 10–13); 63’s questions and source audit (Section 7, Appendix A); 64’s verification, questions and conclusion (Sections 9–11); 65’s certificates, interfaces, questions, conclusion, example and reproduction (Sections 9–12, Appendices A–B); 66’s algorithms, questions, conclusions, the inverse-Pollard relation, dependencies and audit (Sections 12–14, Appendices A–C); 67’s questions and verification (Sections 7–8); 68’s consequences and questions, and verification (Sections 8–9); 69’s certificates, interfaces, verification, questions, conclusion and reproduction (Sections 10–14, Appendix A).

Reading conventions for the sections of sources 60–69 in this Part.

Symbol	Meanings in this Part, by source (<i>tempting false reading</i> in italics)
pins (60–69)	60: 89c1cbc38; 61: 5ea1d06dd; 62: 17123532a; 63: d7b83fa1f; 65: 8c40adc24; 68, 69: 0d93c1eac; 64, 66, 67: none (blobs). Changes at HEAD fb9602e55 to the files they cite are recorded in Section V.1.
questions	as delivered: 60: Research questions 1–14 (own counter); 61, 62: Questions 12.1–12.9, 12.1–12.8; 63: Research questions 7.1–7.13; 64: Questions 10.1–10.8; 65: Research questions 1–9 (own counter); 66: Conjectures 13.1, 13.3 and Questions 13.2, 13.4–13.10; 67: Questions 7.1–7.11; 68: Questions 8.2–8.13 (8.1 is a proposition); 69: Research questions 1–10 (own counter). Here every one is numbered on this article’s counter (the own-counter ones within their section); in Section V.2 they are “60 Q1” and so on, in delivered order.

V.1 The formal project: ledger crosswalk [write]

Added in the write, 6 October 2026.

Every comparison that a source makes with a corrected statement of the formal project was checked against the ledger and the Lean declarations at HEAD 25df8755d (6 October 2026). File and line numbers below are at that commit. Statements are `Sections*.lean` propositions; “exact” means the ledger lists an exact companion theorem. The ledger then records 98 exact companions and 22 open statements; at all three pins it recorded 88 and 32. The last column gives the formal status of the *improvement*. All of the comparisons are accurate at HEAD, except the stale ones the table names.

Gowers	Lean statement; companion (status)	Source claim here	Improvement
Lemma 2.3	<code>lemma_2_3</code> (Sections01_03:141); <code>lemma_2_3_holds</code> (Proofs02Partition:561) (exact)	constant $1/4 \rightarrow 1/3$, optimal under the Lean hypotheses (05, Thm. I.16.4)	not formalized

Gowers	Lean statement; companion (status)	Source claim here	Improvement
Cor. 2.4	<code>corollary_2_4</code> (:158); <code>corollary_2_4_holds</code> (<code>Proofs02PhasePartition</code> :276) (exact); related refutation <code>corollary_2_4_before_roundin</code> <code>g_repair_false</code> (:241) (reordered after the independent check of 6 October 2026; “(exact)” first followed the refutation)	$128\pi \rightarrow 36\pi$, $\alpha/2 \rightarrow 3\alpha/4$, same upper length, $r \leq N$ dropped (05, Cor. I.16.6)	not formalized
Cor. 2.5	<code>corollary_2_5</code> (:180); <code>corollary_2_5_holds</code> (<code>Proofs02</code> <code>DensityIncrement</code> :517); proof-level pair $\sqrt{\alpha N/128\pi}$, $\alpha/4$ inside the proof of private <code>cor25_large_scale</code> (:375; lines 446–509) (exact) (corrected after the independent check of 6 October 2026; first “in private”)	density-sensitive increments (01, 03, 05); 03 and 05 compare with the proof-level pair, 01 with the statement (note after Table 3)	not formalized
Thm. 3.2, Cor. 3.3	<code>theorem_3_2_holds</code> (<code>Proofs03Pr</code> <code>ogressionAverage</code> :381); <code>corollary_3_3</code> (:276), <code>corollary_3_3_holds</code> (<code>Proofs03ProgressionCount</code> :349) (exact; N prime, $k \leq N$)	$2^k \alpha \rightarrow (k-2)\alpha$ in normalized terms; admissible slopes (03, 04, Part III)	not formalized
Cor. 3.6	<code>corollary_3_6</code> (:310); <code>corollary_3_6_holds</code> (<code>Proofs03</code> <code>ProgressionExistence</code> :157) (exact)	density exponent $k2^{k-1}$ against $k2^k$ (04, Cor. III.4.4)	not formalized
Lemmas 3.8, 3.9	<code>lemma_3_8_holds</code> (<code>Proofs03GowersCS</code> :513), <code>lemma_3_9_holds</code> (<code>Proofs03Minkowski</code> :173); tools <code>gowers_cauchy_schwarz</code> (<code>Proofs03GowersCS</code> :485) and <code>gowersNorm_add_le</code> (<code>Proofs03Minkowski</code> :131) (exact)	used as tools (03, 04)	— (corrected after the independent check of 6 October 2026: the row first read “Lemma 3.8” and listed <code>gowersNorm_add_le</code> , the tool of Lemma 3.9, without <code>lemma_3_9_holds</code> ; lines unchanged at <code>e1db00e99</code>)
Lemma 3.10	<code>lemma_3_10</code> (:340); <code>lemma_3_10_holds</code> (<code>Proofs03CubeUpper</code> :82) (exact)	excess $\alpha^{1/2^d} \rightarrow \alpha^{4/2^d}$ (03); leading $6\sqrt{2}\delta^4 u^4$ for $d=3$ (04)	not formalized
Cor. 5.6	<code>corollary_5_6</code> (Section05:205); <code>corollary_5_6_holds</code> (<code>Proofs05FullPartition</code> :85) (exact at HEAD , open at every pin)	02 imports it as $\text{PP}(K, R_0)$ with the same constants; 02’s “ <i>not Lean-verified</i> ” is <i>stale</i> (note in Section I.11)	02’s Thm. I.11.1 not formalized

Gowers	Lean statement; companion (status)	Source claim here	Improvement
Cor. 5.8	<code>corollary_5_8</code> (Section05:242) (open); <code>corollary_5_8_with_scale_holds</code> (Proofs05FullPartition:94) proves a scale-qualified repair	02 declines to restate it (Section I.10)	—
Lemma 5.15	<code>lemma_5_15</code> (Section05:352); <code>lemma_5_15_holds</code> (Proofs05_10:735) (exact; $\alpha \geq 0$, empty cells allowed)	size $\alpha N/(4M) \rightarrow$ $\alpha(2+\alpha)N/(2(4-\alpha)(M-1))$ for $\alpha > 0$, nonempty cells (02, 03)	not formalized
Prop. 6.1	<code>proposition_6_1</code> (Sections06_07:37); <code>proposition_6_1_holds</code> (Proofs05_10:505) (exact)	recovered by 02's Thm. II.2.1 with $m(f) \leq 1$, $\mathcal{U} \leq m^4$	not formalized
Prop. 7.3	<code>proposition_7_3</code> (:135); <code>proposition_7_3_holds</code> (Proofs07BalogSzemerédi:996); also <code>proposition_7_3_group_holds</code> (exact)	$(2^{-20}c^{12}, 2^{38}c^{-24}) \rightarrow$ $(c/8, 2^{19}c^{-9})$ (03)	not formalized
Lemma 7.4	<code>lemma_7_4</code> (:167); <code>lemma_7_4_holds</code> (Proofs07DRC:536), five witnesses (exact)	$2^{-1/2}\delta^5 n \rightarrow 7\delta^3 n/(3\sqrt{6})$ (03)	not formalized
Lemma 9.2	<code>lemma_9_2</code> (Sections08_09:69); <code>lemma_9_2_holds</code> (Proofs09Moments:527); exact $N^2 B $ at line 495 (exact)	(M4) exact Parseval is proof-level; (M5) $\gamma^7\beta^{15}p^{15}$ (04)	not formalized
Lemma 9.3, Cor. 9.4	<code>lemma_9_3</code> (:77), <code>corollary_9_4</code> (:91); companions Proofs09Restriction:376, :416 (exact; <i>no primality</i>)	prime-field restriction, $2^{-95}\gamma^{112}\eta^{15} B ^{15}$ (04); does not prove the Lean statements, as 04 says	not formalized
Lemma 12.4	<code>lemma_12_4</code> (Sections12_13:158); <code>lemma_12_4_holds</code> (Proofs12Amplification:21) (exact)	$\beta^{112} \rightarrow \beta^{106}$ (04, (A12))	not formalized
Cor. 14.6, Lemmas 14.7, 14.8	<code>corollary_14_6</code> (Sections14_15:266), <code>lemma_14_7</code> (:279), <code>lemma_14_8</code> (:292); companions Proofs14Arrangements:838, Proofs14HigherArrangements:712, Proofs14ProductToArrangement s:18 (exact)	profile σ ; $7 \cdot 4^{k+1} \rightarrow 7 \cdot 4^{k+1} - 6$ (04, Thm. II.10.1)	not formalized
Lemma 17.4, Cors. 17.5, 17.6	<code>lemma_17_4_holds</code> (Proofs17Regions:691); <code>corollary_17_5</code> (Sections17_18:103), <code>corollary_17_6</code> (:117); companions Proofs17FinitePatterns:396, :501 (exact)	$2^{2rk^3} \rightarrow (2r)^k C_k$; $M2^{2r(k+1)^3} \rightarrow$ $M \min(2r, M)^k R_{k+1}(k2^{k-1})$ (01, 04, Part IV)	not formalized

Gowers	Lean statement; companion (status)	Source claim here	Improvement
Prop. 17.7	<code>proposition_17_7 (:140);</code> <code>proposition_17_7_holds</code> (<code>Proofs17LocalizedPhaseRemoval:2900</code>), joint	$2^{-2(k+1)^3} \rightarrow 2^{-k^2 - k \log_2 k - O(k)}$ (04); weaker interface (01)	not formalized
Cor. 16.11, Thms. 1.3, 18.1, 18.2	<code>hyperplaneRegionCount</code> (exact) all open ; <code>theorem_18_1</code> 's docstring warns about the tower estimate of Cor. 16.11	01's dependency warning (Section V.3)	—
Section 18 (interval transfer)	<code>hasNatAP_of_short_modular_sequence</code> (<code>Proofs18IntervalTransfer:22</code>), <code>hasNatAP_of_hasModAP_image</code> (:92), added in 327773149	04's Lemma III.4.5; 03's Remark III.1.4	formalized (same statement, $n - 1 \mapsto n$)

Normalizations. The Lean development uses the unnormalized Fourier transform and the unnormalized Gowers norm; the intake found the sources' conversion tables accurate (source 02's Section V.8 and Appendix, source 03's Section V.13); the write did not re-derive them. `Definitions.lean` has changed since the pins only by an added `Box.IsProper`, which no source uses; every other file that a source cites is unchanged.

Formal status of this report. As stated in Section F.5: the lifting lemma is formalized and was formalized before the sources arrived; nothing else here is. The reciprocal pointer for `Combinatorics/Ramsey/FORMALIZATION_STATUS.txt` is left to the session that maintains the Lean development.

[write] **6 October 2026; sources 06–10.** *The ledger at the second write.* At HEAD 764740f07 (6 October 2026) the ledger records 99 exact companions and 21 open statements, and `FORMALIZATION_STATUS.txt` reports 196 modules reachable from the audit and 1,805 audited public theorems. Since the first write's HEAD 25df8755d the statements closed are: `lemma_16_9`. Since 7d99e5dee, the pin of sources 06, 09 and 10 (95 exact, 25 open, 140 modules, 1,365 theorems), 4 statements have closed: `lemma_16_1`, `lemma_16_5`, `lemma_16_6` and `lemma_16_9`; since ae0e89574, the pin of source 07, 4: `lemma_16_1`, `lemma_16_5`, `lemma_16_6` and `lemma_16_9`. The table above keeps its file and line numbers at 25df8755d; the rows below give them at 764740f07. The Lean development is being extended by another session as this is written, so the counts may already have moved. Every comparison made by sources 06–10 with a formal statement was checked against these declarations and is accurate, except the stale counts corrected in the notes after source 06's Section F.11 and source 10's Section F.15.

Gowers	Lean statement; companion (status)	Claim of sources 06–10	Improvement
Cor. 2.5	<code>corollary_2_5</code> (<code>Sections01_03:180</code>); <code>corollary_2_5_holds</code> (<code>Proofs02DensityIncrement:517</code>) (exact)	increment $\delta\eta/(4\delta - \eta)$ at source 01's length; the frontier D_γ, L_γ (10, Thm. I.24.2)	not formalized
Thm. 2.6	<code>theorem_2_6 (:189);</code> <code>theorem_2_6_holds</code> (<code>Proofs02Roth:1238</code>) (exact)	step count $321(\delta_0^{-1} - 1)$ against the transcription's $640\delta_0^{-1}$ (10, Prop. V.42.1)	not formalized

Gowers	Lean statement; companion (status)	Claim of sources 06–10	Improvement
Thm. 3.2, Cors. 3.3, 3.6	<code>theorem_3_2_holds</code> (Proofs03ProgressionAverage:381); <code>corollary_3_3</code> (:276), <code>corollary_3_3_holds</code> (Proofs03ProgressionCount:349); <code>corollary_3_6</code> (:310), <code>corollary_3_6_holds</code> (Proofs03ProgressionExistence:157) (exact)	two-target bound and quadratic discrepancy, $\alpha \leq (\delta^{k-1/2}/2B_k)^{2^{k-2}}$ (06); torsion kernels, error $\delta S_k \eta$, $\alpha \leq (\delta^{k-1}/2S_k)^{2^{k-1}}$ (10)	not formalized
Lemmas 3.7, 3.10	<code>lemma_3_7</code> (:320), <code>lemma_3_7_holds</code> (Proofs03Cubes:184); <code>lemma_3_10</code> (:340), <code>lemma_3_10_holds</code> (Proofs03CubeUpper:82) (exact)	excess of fourth order with exact quartic term (08, 09) and quintic term (09); exponent $4/2^d$ sharp for sets; lower bound $(\delta^4 + E_4)^{2^d/4}$ (08)	not formalized
Lemma 5.15	<code>lemma_5_15</code> (Section05:352); <code>lemma_5_15_holds</code> (Proofs05_10:735) (exact)	fixed- M optimum, $\tau(\alpha)$ (06); cap and tail (10)	not formalized
Prop. 7.3	<code>proposition_7_3</code> (Sections06_07:135); <code>proposition_7_3_holds</code> (Proofs07BalogSzemerédi:996) (exact)	$(c/4, 2^{10}c^{-5})$ on any abelian group (10)	not formalized
Lemma 7.5, Cor. 7.6	<code>lemma_7_5</code> (:183), <code>lemma_7_5_holds</code> (Proofs07FreimanRestriction:267); <code>corollary_7_6</code> (:197), <code>corollary_7_6_holds</code> (Proofs07AdditiveRestriction:100) (exact)	$ B /(k(h_k - 1) + 1)$; $2^{-209}\gamma^{103}$ against $2^{-1882}\gamma^{1164}$ (10); both imply the formal conclusions (note after Cor. II.21.2)	not formalized
Lemma 10.9	<code>lemma_10_9</code> (Section10:335); <code>lemma_10_9_holds</code> (Proofs10Shift:717) (exact; FreimanHom 2)	order five under exactly these hypotheses; order nine with the hypotheses of <code>lemma_10_6</code> and the set-up (07)	not formalized; order nine needs a new statement
Cors. 17.5, 17.6	as in the table above (Sections17_18:103, :117; Proofs17FinitePatterns:396, :501) (exact)	$(2r)^k C_k$, $\min(2r, M)^k C_{k,M}$, entropy uniform in M ; $C_4 = 154$ certified (06)	not formalized
Prop. 17.7	<code>proposition_17_7</code> (:140); <code>proposition_17_7_holds</code> (Proofs17LocalizedPhaseRemoval:2900) (exact)	$1/(8K_k^2) = 2^{-2k^2 - O(k \log k)}$ (06)	not formalized

Gowers	Lean statement; companion (status)	Claim of sources 06–10	Improvement
Cor. 16.11, Thms. 18.1, 18.2	<code>corollary_16_11</code> (Section16:893), <code>theorem_18_1</code> (Sections17_18:171), <code>theorem_18_2 (:190): open;</code> <code>section18_cyclic_density_incre</code> <code>ment_holds_of_theorem_18_1</code> (Proofs18QuantitativeSzemeredy:169)	dependency warnings (06, 10); iteration potential (06)	—

[write] 6 October 2026; sources 06–10. *Formal status after the second write.* No statement of sources 06–10 is formalized, and none gains formal status by its placement. The only formalized statement of the report is still the lifting lemma of source 04. Sources 06–10 ship no Lean files; their formalization plans are in Part V (source 07’s also as 07-transport-rigidity-FORMALIZATION.md).

[write] 6 October 2026; sources 11–14. *The ledger at the third write.* At HEAD `e1db00e99` (6 October 2026) the ledger records 100 exact companions and 20 open statements, and `FORMALIZATION_STATUS.txt` reports 206 modules and 1,879 audited public theorems; since the second write’s HEAD `764740f07` the statement closed is `lemma_13_8`. At `40f885098` (source 11’s pin) and `9ea383e9b` (the pin of sources 12–14) it recorded 97 and 23, 154 modules and 1,508 theorems; the statements closed since are `lemma_16_6`, `lemma_16_9` and `lemma_13_8`. The rows below give file and line numbers at `e1db00e99`. Every comparison of sources 11–14 with a formal statement was checked against these declarations and is accurate, except the stale counts of sources 12 and 13 and source 12’s incomplete list of open obligations (notes in Sections F.17, F.18 and II.41), and two location slips (source 14 places `HasMonochromaticAP` in `Sections08_09.lean`, where only `corollary_8_3` is; source 13’s interface table attaches its envelope to `Proofs05Downstream.lean`, while the companion of Lemma 5.14 is `lemma_5_14_holds`, `Proofs05FullPartition.lean`, line 103, via `lemma_5_14_holds_of_corollary_5_6`, `Proofs05Lemma14.lean`, line 231).

Gowers	Lean statement; companion (status)	Claim of sources 11–14	Improvement
Cor. 2.5	<code>corollary_2_5</code> (Sections01_03:180); <code>corollary_2_5_holds</code> (Proofs02DensityIncrement:517) (exact)	increment $\delta + \alpha / (4 - \alpha(2 - \delta) / (\delta(1 - \delta)))$ at source 01’s length (13); dominated by source 01’s (note after Thm. I.29.1)	not formalized
Thm. 3.2, Cor. 3.3	<code>theorem_3_2_holds</code> (Proofs03ProgressionAverage:381); <code>corollary_3_3</code> (Sections01_03:276), <code>corollary_3_3_holds</code> (Proofs03ProgressionCount:349) (exact)	telescope with error $\sum_j \ f_j\ _{U^{j-1}} \prod_{i>j} \delta_i$ (13; as 03, 04)	not formalized
Lemma 3.5	<code>lemma_3_5</code> (Sections01_03:299); <code>lemma_3_5_holds</code> (Proofs03Interval:412) (exact; constant 3)	$2\sqrt{2}\alpha^{1/4}N \min(\beta, 1 - \beta)^{1/4}$, exact dual constant (13)	not formalized

Gowers	Lean statement; companion (status)	Claim of sources 11–14	Improvement
Lemma 3.10	<code>lemma_3_10</code> (Sections01_03:340); <code>lemma_3_10_holds</code> (Proofs03CubeUpper:82); <code>cubeCount</code> (Definitions:215) (exact)	three-cube orbit identity (13); secondary term $(9/4)^{1/3}\delta^{8/3}\alpha^{2/3}$, $\gcd(N, 6) = 1$, small α (14)	not formalized
Lemma 5.9	<code>lemma_5_9</code> (Section05:261); <code>lemma_5_9_holds</code> (Proofs05FullPartition:98) (exact); one-phase <code>section5_strong_linear_partition</code> (Proofs05LinearPartition:86)	linear case with q phases, exponent $1/(3q + 2)$, no threshold (13)	not formalized
Lemmas 5.14, 5.15	<code>lemma_5_14_holds</code> (Proofs05FullPartition:103); <code>lemma_5_15</code> (Section05:352), <code>lemma_5_15_holds</code> (Proofs05_10:735) (exact)	sharp weighted selection and joint frontier (13; as 01, 02, 03, 06)	not formalized
Prop. 6.1, Lemma 6.2	<code>proposition_6_1</code> (Sections06_07:37), <code>proposition_6_1_holds</code> (Proofs05_10:505); <code>lemma_6_2</code> (:47), <code>lemma_6_2_holds</code> (Proofs05_10:589); <code>functionGraph</code> (:29) (exact)	affine recovery for $\alpha^4 > 8/9$; exact endpoint profiles; the formal statement needs no primalty (11)	not formalized
Cor. 8.3	<code>corollary_8_3</code> (Sections08_09:52) open ; <code>corollary_8_3_holds_of_theorem_8_2</code> (Proofs08Coloring:307); <code>HasMonochromaticAP</code> (Definitions:113)	a different, symmetric predicate over $\mathbb{F}_q^d$ (14, as 10)	—
Lemmas 10.3–10.9	<code>lemma_10_3_holds</code> (Proofs10Selection:106), <code>lemma_10_5_holds</code> (Proofs10Regular:782), <code>lemma_10_6_holds</code> (Proofs10Extraction:1836, with <code>extractionMode</code> , :1532), <code>lemma_10_7_holds</code> (Proofs05_10:954), <code>lemma_10_8_holds</code> (Proofs10Shift:480), <code>lemma_10_9_holds</code> (:717) (exact)	12’s component input (Lemma II.40.1) and its Appendix A repair; borrowed, not new	formal input
Lemma 10.10, Cor. 10.11, Lemma 10.12	<code>lemma_10_10</code> (Section10:362), <code>lemma_10_10_holds</code> (Proofs10Bohr:277); <code>corollary_10_11</code> (:374), <code>corollary_10_11_holds</code> (Proofs10Bohr:607); <code>lemma_10_12</code> (:401), <code>lemma_10_12_holds</code> (Proofs10Induced:145) (exact)	sharp extension $3\varepsilon + 2\kappa < 1$; radius $\rho/(8d)$ after selection (an alternative, not a drop-in replacement); synchronization (12)	not formalized

Gowers	Lean statement; companion (status)	Claim of sources 11–14	Improvement
Thm. 10.13	<code>theorem_10_13</code> (Section10:473); <code>theorem_10_13_holds</code> (Proofs10Main:1292) (exact; $\alpha \leq 1/6$, $\alpha^6 X /20000$); proof-level $ W \geq \varrho^2 \alpha^2 LN/16$ (:1398), $\alpha \leq 1/6$ spent at :1407	radius $\alpha^{19}/(3\pi 2^{149})$; fraction $\alpha^5/20000$, already proof-level as $5\alpha^5/98304$; all densities (12)	radius not formalized
Cor. 10.14	<code>corollary_10_14</code> (Section10:501); <code>corollary_10_14_holds</code> (Proofs 10ProgressionLinearity:133); <code>HasBohrDifferenceModel</code> (:458) (exact)	scalar linearity on the larger Bohr set (12)	follows from the formal corollary given a model
Lemmas 12.3, 14.7, 14.8	<code>lemma_12_3</code> (Sections12_13:151), <code>lemma_12_3_holds</code> (Proofs12Hi gherArrangements:334); <code>lemma_14_7</code> (Sections14_15:279), <code>lemma_14_7_holds</code> (Proofs14Hi gherArrangements:712); <code>lemma_14_8_holds</code> (Proofs14Pr oductToArrangements:18); <code>respectedCubePairCount</code> (Sections14_15:183) (exact)	equality and stability of the profile bound; denominators $S - \Xi$, S_{card} (14)	not formalized
Section 18 (interval transfer)	<code>hasNatAP_of_short_modular_se quence</code> (Proofs18IntervalTransfer:22)	13's remark that its modular criterion needs this transfer	formalized (before arrival)

[write] **6 October 2026; sources 11–14.** *Formal status after the third write.* No statement of sources 11–14 is formalized, and none gains formal status by its placement. The only formalized statement of the report is still the lifting lemma of source 04. Sources 11–14 ship no Lean files; their formalization plans are in Part V (source 11's also as `11-affine-rigidity-formalization_plan.md`).

[write] **6 October 2026; sources 15–20.** *The ledger at the fourth write.* At HEAD `0b4890a1c` (6 October 2026) the ledger records 101 exact companions and 19 open statements, and `FORMALIZATION_STATUS.txt` reports 233 modules and 2,138 audited public theorems; since the third write's HEAD `e1db00e99` the statement closed is `lemma_13_7` (`lemma_13_7_holds`, `Proofs13CompleteRowExtraction.lean`, line 46). At the pins of sources 15–20 it recorded: `9ea383e9b` (source 15) 97 and 23, 154 modules, 1,508 theorems; `dca9f0638` (sources 16–19) 98 and 22, 169 modules, 1,658 theorems; `1bd57e91c` (source 20's checkpoint) 99 and 21, 173 modules, 1,689 theorems. Closed since: `lemma_16_6` (since `9ea383e9b` only), `lemma_16_9` (since `dca9f0638`), `lemma_13_8` and `lemma_13_7`. The rows below give file and line numbers at `0b4890a1c`; every one was checked by reading the declaration at that line. Every comparison of sources 15–20 with a formal statement was checked against these declarations and is accurate, except: the stale counts of sources 16 and 19, source 19's status of Lemma 16.9 and its remark on missing properness clauses (notes in Sections V.65 and V.80); and three understatements of formal status, recorded in dated notes: the unweakened Corollary 14.6 exponent is kernel-checked as a private lemma (source 15), Proposition 17.7 is an exact companion (source 16), and the all- N Lemma 9.3 and Corollary 9.4 are exact companions (source 17). The file and line citations of the three earlier tables were re-checked at `0b4890a1c` as well (dated note after the last table).

Gowers	Lean statement; companion (status)	Claim of sources 15–20	Improvement
Thm. 3.2, Cor. 3.3	<code>theorem_3_2</code> (Sections01_03:258), <code>theorem_3_2_holds</code> (Proofs03Pr ogressionAverage:381); <code>corollary_3_3</code> (:276), <code>corollary_3_3_holds</code> (Proofs03ProgressionCount:349) (exact; N prime, $k \leq N$)	two marked functions with an L^2 gain for surjective homomorphisms; quadratic counting $B_k(\delta)u^2$, strictly below source 06’s coefficient (20)	not formalized
Lemma 3.10	<code>lemma_3_10</code> (Sections01_03:340), <code>lemma_3_10_holds</code> (Proofs03CubeUpper:82) (exact)	expansion through order five on every finite abelian group (16, 18; as 09)	not formalized
Lemmas 5.5, 5.9	<code>lemma_5_5</code> (Section05:181), <code>lemma_5_5_holds</code> (Proofs05Poly nomialPartition:1678); <code>lemma_5_9</code> (:261), <code>lemma_5_9_holds</code> (Proofs05FullPartition:98) (exact)	exponents polynomial in the number of phases, conditional on Maynard’s theorem (18, 19); shared leading coefficient through the scalar Lemma 5.5 (18)	not formalized
Lemmas 9.2, 9.3, Cor. 9.4	<code>lemma_9_2</code> (Sections08_09:69), <code>lemma_9_2_holds</code> (Proofs09Moments:527); <code>lemma_9_3</code> (:77), <code>corollary_9_4</code> (:91); companions <code>Proofs09Restriction:376, :416</code> (exact; all N); private <code>section9_second_moment_le</code> (Proofs09Moments:480)	moments without density loss (17, as 04); prime-modulus restriction with $\alpha^{16}\eta^{15}/(2 \cdot 18^{16})$ and an explicit N_0 , which implies the prime case of <code>lemma_9_3</code> and <code>corollary_9_4</code> (17)	not formalized
Lemma 13.5	<code>lemma_13_5</code> (Sections12_13:410) <code>open</code> ; <code>IsStage135Data</code> (:397); <code>le mma_13_5_of_candidate_family</code> (Proofs13QuadraticRecurrenc e:425); proved at large prime N : <code>lemma_13_5_large_N</code> (Proofs13 LargeScaleRecurrence:96)	own-step partitions with $\gamma < 1/((q+1)(42q+1))$, conditional on Lau’s preprint; weighted cell $\alpha^{32}/16$ (16)	—; Lau would have to be an axiom
Lemma 14.4, Cor. 14.5	<code>lemma_14_4_holds</code> , <code>corollary_14_5_holds</code> (Proofs14Configurations:698, :942) (exact)	15’s configuration and collision lower bounds (as 15 says)	formal input
Cor. 14.6, Lemmas 14.7, 14.8	<code>corollary_14_6</code> (Sections14_15:266), <code>lemma_14_7</code> (:279), <code>lemma_14_8</code> (:292); companions <code>Proofs14Arrangements:838</code> , <code>Proo fs14HigherArrangements:712</code> , <code>Proofs14ProductToArrangement s:18</code> (exact); private <code>cor146_direct_lower</code> (Proofs14Arrangements:759), <code>cor147_moment_eq_count</code> (Proo fs14HigherArrangements:484), <code>cor147_interpolation</code> (:653)	coupled $\beta^{22 \cdot 4^k} \gamma^{44k4^k + 56 \cdot 2^k}$ against $\beta^{28 \cdot 4^k} \gamma^{84k4^k}$ (15)	not formalized

Gowers	Lean statement; companion (status)	Claim of sources 15–20	Improvement
Lemmas 15.4, 15.5	lemma_15_4 (Sections14_15:402), lemma_15_4_holds (Proofs15DegenerateArrangements:1131); lemma_15_5 (:414), lemma_15_5_holds (Proofs15Restriction:3086) with the constant $3^{16 \cdot 2^k} k$ of private selectionDegenerate_real_small (:2690) (exact)	sharp density $(k + C_d)/p$, below the formal bound for $d \geq 2$, $p > 2d$; $\lambda p \geq k + 2,585,445$ (18)	not formalized
Lemmas 16.9, 16.10	lemma_16_9 (Section16:749), lemma_16_9_holds (Proofs16Lemma9:158) (exact since 16bbe5ad); lemma_16_10 (:851) open ; Section16AllBoxLineCovers (:824, proper boxes)	common sampling, shared-slice lifting and a local replacement; Lemma 16.10 itself not proved (19)	not formalized
Section 16, formal results after the pins	section16_printed_distinct_sampling_counterexample (Proofs16DistinctSampling:204), section16_distinct_sampling_budget (:117), section16_same_dimension_closing_comparison_fails (Proofs16ClosingComparison:25), section16_local_affine_lift (Proofs16LocalAffineLift:13), section16_local_affine_lift_all_scales (Proofs16LiftAllScales:46), section16_local_affine_lift_sqrt (Proofs16LiftWidth:31)	confirm 19’s audit points; weaker analogues of 19’s Corollary II.70.8	formal (not 19’s statements)
Prop. 17.7	proposition_17_7 (Sections17_18:140), UniformOnPartition (:132); proposition_17_7_holds (Proofs17LocalizedPhaseRemoval:2900) (exact); private prop177_small_branch (:1283), powersetDifference_pow, polarizationPhase (Proofs17Phase:142, :217)	$(4/(3(k+2)))^{k+2}$ in the same normalization: implies the formal statement (16)	not formalized
Cors. 17.5, 17.6; Thm. 18.1	corollary_17_5_holds, corollary_17_6_holds (Proofs17FinitePatterns:396, :501) (exact); theorem_18_1 (Sections17_18:171) open	the carry counts remain needed for the corollaries; larger cells must be transported in Section 18 (16)	—

[write] 6 October 2026; sources 15–20. *Formal status after the fourth write.* No statement of sources 15–20 is formalized, and none gains formal status by its placement. The only formalized statement of the report is still the lifting lemma of source 04. Four statements imply the conclusions of formal statements with

better constants, none of them formally: source 16’s Theorem F.21.1 implies `proposition_17_7` in all cases; source 17’s Theorem F.22.1 and its order-eight corollary imply `lemma_9_3` and `corollary_9_4` for prime N ; source 18’s Theorem II.66.8 implies the bound of `lemma_15_4` for $d \geq 2$ and primes $p > 2d$; source 18’s Corollary II.66.9 would sharpen an internal constant of `lemma_15_5_holds`. Sources 15–20 ship no Lean files; their formalization plans are in Part V (source 15’s also as `15-collision-amplification-integration.md`).

Re-check of the earlier tables at 0b4890a1c. The file-and-line references of the three earlier tables (first write at 25df8755d, second at 764740f07, third at e1db00e99; 211 references counted with repetitions) were re-read at 0b4890a1c by a script that opens each cited file at the cited line and reads the declaration that begins there. All are correct at 0b4890a1c; seven of them name their file only through an earlier row and were resolved by hand. The counts printed with each table are those of its own HEAD.

[write] 6 October 2026; sources 21–27. *The ledger at the fifth write.* At HEAD 2dfe46f44 (6 October 2026) the ledger records 101 exact companions and 19 open statements, and `FORMALIZATION_STATUS.txt` reports 267 modules and 2,424 audited public theorems; no statement has closed since the fourth write’s HEAD 0b4890a1c. At the pins of sources 21–26 it recorded: 43306653a (sources 21 and 22) 99 and 21, 175 modules, 1,716 theorems; 65f4e039f (source 23) 99 and 21, 177 modules, 1,734 theorems; 2e0553b0c (source 24) 97 and 23, 157 modules, 1,547 theorems; 568d1802e (source 25) 99 and 21, 180 modules, 1,744 theorems; c94998a5d (source 26) 99 and 21, 185 modules, 1,760 theorems. Closed since: `lemma_13_8` and `lemma_13_7` (and, since 2e0553b0c only, `lemma_16_6` and `lemma_16_9`). None of sources 21–27 quotes ledger counts, and every Lean file they cite has the same blob at HEAD as at their pins (source 27: the blobs it records), so none of their repository statements is stale. The rows below give file and line numbers at 2dfe46f44; every one was checked by reading the declaration at that line. Every comparison of sources 21–27 with a formal statement was checked against these declarations and is accurate, except one gap overstated by source 24 (the abelian-group form of Proposition 7.3 is already formal; note in Section F.29). The row of the first table on Corollary 5.8 (“02 declines to restate it”) stands; source 26 now supplies the degree-one case (second row below).

Gowers	Lean statement; companion (status)	Claim of sources 21–27	Improvement
Lemma 3.10	<code>lemma_3_10</code> (Sections01_03:340), <code>lemma_3_10_holds</code> (Proofs03CubeUpper:82) (exact)	exponent five: the sharp coefficient $2^{3/4}/3$ with an absolute $O(u^8)$ remainder; 20/9 next on $\mathbb{Z}/5\mathbb{Z}$ (21; also 23)	not formalized
Cor. 5.8	<code>corollary_5_8</code> (Section05:242) open ; <code>corollary_5_8_with_scale_holds</code> (Proofs05FullPartition:94); <code>PolynomialOn</code> (Definitions:222), <code>polynomialPartitionConstant</code> (Section05:45)	degree one without a scale threshold: gain $\alpha/3$ at length $(N/M)^{1/3}/8$; implies the $k = 1$ instance, which is trivial for $N/M \leq 2^{48}$ (26; write)	not formalized; candidate (ledger unchanged)
Lemma 5.9	<code>lemma_5_9</code> (Section05:261), <code>lemma_5_9_holds</code> (Proofs05FullPartition:98) (exact)	prescribed-length thresholds with matching lower bounds; source 13’s linear exponent $1/(3q + 2)$ is optimal (23)	not formalized

Gowers	Lean statement; companion (status)	Claim of sources 21–27	Improvement
Prop. 7.3, Lemmas 7.4, 7.5, Cor. 7.6	<code>proposition_7_3</code> (Sections06_07:135), <code>proposition_7_3_holds</code> (Proofs07BalogSzemerédi:996), <code>proposition_7_3_group_holds</code> (:854); <code>lemma_7_4</code> (Sections06_07:167), <code>lemma_7_4_holds</code> (Proofs07DRC:536); <code>lemma_7_5</code> (Sections06_07:183), <code>lemma_7_5_holds</code> (Proofs07FreimanRestriction:267); <code>corollary_7_6</code> (Sections06_07:197), <code>corollary_7_6_holds</code> (Proofs07AdditiveRestriction:100) (exact)	second routes to sources 03 and 10 with equal or weaker constants; extraction for $\mathbb{F}_p^d$ codomains (24)	not formalized; the group form is already formal
Lemmas 15.4, 15.5	<code>lemma_15_4</code> (Sections14_15:402), <code>lemma_15_4_holds</code> (Proofs15DegenerateArrangements:1131); <code>lemma_15_5</code> (Sections14_15:414), <code>lemma_15_5_holds</code> (Proofs15Restriction:3086); <code>arrangementSelectionExponent</code> (Sections14_15:409), <code>GeneralArrangement</code> (:192), <code>degenerateGeneralArrangementCount</code> (:364) (exact)	$k + C_d$ again (25, 27); $d = 1$ (22); second coefficient for $d = 2$ (27); stratified finite bound (22); selection for every d (25); all below the formal bound in their ranges	not formalized
Lemma 16.10	<code>lemma_16_10</code> (Section16:851) open ; <code>Box.short_parent_partition</code> (Proofs16ShortParents:23), <code>box_product_tiling_of_contained_axis</code> (Proofs16RelativeSteps:89), <code>section16_local_affine_lift_sqrt</code> (Proofs16LiftWidth:31), <code>section16_distinct_sampling_budget</code> (Proofs16DistinctSampling:117)	a common base partition with width $(m/3)^A$ and exact transport, without inherited coordinates; a sampling lemma (23)	not formalized; candidate

[write] 6 October 2026; sources 21–27. *Formal status after the fifth write.* No statement of sources 21–27 is formalized, and none gains formal status by its placement. The only formalized statement of the report is still the lifting lemma of source 04. Two formalization candidates are recorded, neither claimed: the degree-one instance of `corollary_5_8` (source 26; trivial below $N/M = 2^{48}$) and source 23’s common base partition. Sources 22, 25 and 27 imply the bound of `lemma_15_4` in their ranges, like source 18. Sources 21–27 ship no Lean files; source 27’s `FORMALIZATION.md` (shipped as `27-degeneracy-second-order-FORMALIZATION.md`) is a plan.

Re-check of the earlier tables at 2dfe46f44. The file-and-line references of the four earlier tables (first write at 25df8755d, second at 764740f07, third at e1db00e99, fourth at 0b4890a1c) were re-read at 2dfe46f44

by the same script; all are correct at 2dfe46f44, those resolved by hand at the fourth write included, except one: `proposition_17_7_holds`, printed at line 2900 of `Proofs17LocalizedPhaseRemoval.lean` in the first, second and fourth tables (and in some earlier dated notes), is at line 2941 since c2975a00d, which inserted `proposition_17_7_with_upper` (line 2911) above it. Line 2900 was right at every earlier write’s HEAD, including bf2fe146d; the dated tables keep their numbers. The counts printed with each table are those of its own HEAD.

[write] **6 October 2026; sources 28–33.** *The ledger at the sixth write.* At HEAD 5d4718772 (6 October 2026) the ledger records 105 exact companions and 15 open statements, and `FORMALIZATION_STATUS.txt` reports 297 modules and 2,676 audited public theorems; since the fifth write’s HEAD 2dfe46f44 (101 and 19) `theorem_7_1` and `theorem_7_2` (56e774055) and `theorem_8_2` and `corollary_8_3` (6a6d7961c) have closed. At the pins of sources 28–33 it recorded: c94998a5d (source 32) 99 and 21, 185 modules, 1,760 theorems; e3125e698 (sources 28, 29, 31) and 9f83dbe0f (source 33) 99 and 21, 196 modules, 1,805 theorems; 8d65d3c9f (source 30) 100 and 20, 199 modules, 1,823 theorems. Closed since those pins: `lemma_13_8` (not since 8d65d3c9f), `lemma_13_7` and the four above; none is a statement sources 28–33 compare with. Since 128f514af the first obligation of `FORMALIZATION_STATUS.txt` (Corollary 5.8) names source 41 as a written proof of the whole statement and sources 26 and 32 as its degree-one and degree-two precursors; no formal status changes. None of sources 28–33 quotes ledger counts, and every Lean file they cite has the same blob at HEAD as at their pins, so none of their repository statements is stale. The rows below give file and line numbers at 5d4718772; every one was checked by reading the declaration at that line. Every comparison of sources 28–33 with a formal statement was checked against these declarations and is accurate. The fifth table’s Corollary 5.8 row stands; source 32 now supplies the degree-two case (second row below).

Gowers	Lean statement; companion (status)	Claim of sources 28–33	Improvement
Lemma 3.10	<code>lemma_3_10</code> (Sections01_03:340), <code>lemma_3_10_holds</code> (Proofs03CubeUpper:82) (exact)	$c_4^{\text{odd}} \leq (8/35)^{1/4}$, $c_d^{\text{odd}} \geq 1/3$ (29); $c_4^{\text{odd}} \geq 0.5587294$, no single-cosine extremizer for $d \geq 4$ (28)	not formalized
Cor. 5.8	<code>corollary_5_8</code> (Section05:242) <code>open</code> ; <code>PolynomialOn</code> (Definitions:222), <code>polynomialPartitionConstant</code> (Section05:45), <code>ModAP.IsProper</code> (Definitions:93)	degree two without a scale threshold: gain $\alpha/3$ at length $\max\{1, (N/M)^{1/39}/17\}$; implies the $k = 2$ instance (32; with 26, $k \leq 2$)	not formalized; candidate (ledger unchanged)
Prop. 6.1	<code>proposition_6_1</code> (Sections06_07:37), <code>proposition_6_1_holds</code> (Proofs05_10:505) (exact)	sharp near-extremal U^3 profile and selected-graph stability (31); the endpoint with soft and mixed selectors (29)	not formalized
Lemma 7.5, Cor. 7.6	<code>lemma_7_5</code> (Sections06_07:183), <code>lemma_7_5_holds</code> (Proofs07Frei manRestriction:267); <code>corollary_7_6</code> (Sections06_07:197), <code>corollary_7_6_holds</code> (Proofs07 AdditiveRestriction:100) (exact)	weighted restriction $1/(kC^{2k+1})$ (source 10’s with weights); correlation-preserving restriction on $\mathbb{F}_p$ (31)	not formalized
Lemma 9.3	<code>lemma_9_3</code> (Sections08_09:77), <code>lemma_9_3_holds</code> (Proofs09Restriction:376) (exact)	fixed-field restriction keeping every degenerate tuple, threshold $8qS/(\alpha\eta)$, then none (29)	not formalized; does not imply the all-moduli statement

Gowers	Lean statement; companion (status)	Claim of sources 28–33	Improvement
Lemma 10.10, Cor. 10.11, Lemma 10.12, Thm. 10.13	<code>lemma_10_10</code> (Section10:362), <code>lemma_10_10_holds</code> (Proofs10Bohr:277); <code>corollary_10_11</code> (Section10:374), <code>corollary_10_11_holds</code> (Proofs10Bohr:607); <code>lemma_10_12</code> (Section10:401), <code>lemma_10_12_holds</code> (Proofs10Induced:145); <code>theorem_10_13</code> (Section10:473), <code>theorem_10_13_holds</code> (Proofs10Main:1292) (exact)	fixed-radius escape inequality, radius $\rho/(8(3^r - 1))$, 5/14 (30); local model with $\alpha p(\alpha)/20$, $O(\alpha^{-3})$, $\Omega(\alpha^5)$ (28)	not formalized; source 28 against <code>theorem_10_13</code> is a marked question
Lemma 12.3	<code>lemma_12_3</code> (Sections12_13:151), <code>lemma_12_3_holds</code> (Proofs12HigherArrangements:334) (exact); <code>lemma_12_5</code> (:170), <code>lemma_12_6</code> (:183)	$\max\{\theta^7, 4\theta - 3\}N^{32}$; sharp full-product repair (33)	not formalized; candidate

[write] 6 October 2026; sources 28–33. *Formal status after the sixth write.* No statement of sources 28–33 is formalized, and none gains formal status by its placement. The only formalized statement of the report is still the lifting lemma of source 04. Formalization candidates, none claimed: the instances $k \leq 2$ of `corollary_5_8` (sources 26 and 32) and the linear branch of `lemma_12_3` (source 33). Sources 28–33 ship no Lean files; the `FORMALIZATION.md` files of sources 30 and 33 (shipped as `30-fixed-radius-bohr-FORMALIZATION.md`, `33-arrangement-stability-FORMALIZATION.md`) are plans.

Re-check of the earlier tables at 5d4718772. Since the fifth write’s HEAD 2dfe46f44 the tracked Lean files of the development have changed only by new modules, an extended facade and lines appended at the end of `Proofs05Downstream.lean` (after line 966; earlier text cites only its line 244); no cited line moved. The same script re-read their 350 references at 5d4718772 and found exactly the items recorded at the fifth write (the hand-resolved list-pairing artifacts and the move of `proposition_17_7_holds` to line 2941); the dated tables keep their numbers.

[write] 6 October 2026; sources 34–39. *The ledger at the seventh write.* At HEAD 89c1cbc38 (6 October 2026) the ledger records 107 exact companions and 13 open statements, and `FORMALIZATION_STATUS.txt` reports 410 modules and 2,840 audited public theorems; since the sixth write’s HEAD 5d4718772 `lemma_16_4` (c450a0c73) and `theorem_1_2` (5ea1d06dd, Szemerédi’s theorem from a vendored hypergraph-removal theorem) have closed. At the pins of sources 34–39 it recorded 100 and 20: 12da35673 (sources 34, 35, 38) and 6c2e2172b (source 36) 206 modules and 1,878 theorems, 869289aa3 (sources 37, 39) 206 modules and 1,879 theorems. Closed since those pins: `lemma_13_7`, `theorem_7_1`, `theorem_7_2`, `theorem_8_2`, `corollary_8_3`, `lemma_16_4` and `theorem_1_2`; none is a statement sources 34–39 compare with (source 39 treats the sampling step of Lemma 16.10, not Lemma 16.4). Only source 38 quotes ledger counts (those of its pin; note in its Section V.143); every Lean file and every paper file that sources 34–39 cite has the same blob at HEAD as at their pins (the intake’s provenance check, rerun at the write: 56 records, all matching), so none of their statements about the formal project is stale apart from those counts. The rows below give file and line numbers at 89c1cbc38; every one was checked by reading the declaration at that line, and every comparison of sources 34–39 with a formal statement was checked against these declarations.

Gowers	Lean statement; companion (status)	Claim of sources 34–39	Improvement
Lemma 3.10	lemma_3_10 (Sections01_03:340), lemma_3_10_holds (Proofs03CubeUpper:82) (exact)	$c_4^{\text{odd}} \leq 2/3$, leading cube coefficient $200/3$ at $d = 4$, β_d for $5 \leq d \leq 8$, prime-torsion onset $h(p) = D_{\text{div}}(p) + 1$ (34); coefficient 5 (35); exact quotient reduction in exponent five, u^{10}, u^{12} (38)	not formalized
Lemma 5.2	lemma_5_2 (Section05:141), lemma_5_2_holds (Proofs05MissingInterval:220) (exact)	frequency $\leq \lfloor N/M \rfloor - 1$, coefficient $ A M/(N - M)$, any weights, no parity of M (37)	implies the formal conclusion ([write] note); not formalized; candidate
Lemmas 5.3, 5.5	lemma_5_3 (Section05:152), lemma_5_3_holds (Proofs05Weyl:3607); lemma_5_5 (Section05:181), lemma_5_5_holds (Proofs05PolynomialPartition:1678); lemma_5_5_square_root_auxiliary (Section05:192) (exact)	recurrence exponent $(1 - 1/(4k))/(2^{k-1} + 1)$ from the formal Weyl input, $7/24$ against $1/16$ at $k = 2$ (37)	implies the formal bound for $t \geq \text{weylThreshold}$ ([write] note); not formalized
Lemma 5.9	lemma_5_9 (Section05:261), lemma_5_9_holds (Proofs05FullPartition:98) (exact)	necessary exponents $(W + 2)a + \sum d_i b_i \leq 1$ for prescribed lengths; exact thresholds $DC(S)$ (35)	an obstruction; no formal change
Prop. 6.1	proposition_6_1 (Sections06_07:37), proposition_6_1_holds (Proofs05_10:505) (exact)	sharp odd-order stability modulus $\varepsilon/4 + 3\varepsilon^2/16 + (\sqrt{2}/64)\varepsilon^{5/2}$ (38)	not formalized
Lemma 10.9	lemma_10_9 (Section10:335), lemma_10_9_holds (Proofs10Shift:717) (exact)	sharp nonuniform threshold $1/(2 + (2k - 2)\kappa)$ in source 07's finite-fibre model (34)	a different model; no formal change
Lemmas 14.8, 15.5	lemma_14_8 (Sections14_15:292), lemma_14_8_holds (Proofs14ProductToArrangements:18); lemma_15_5 (Sections14_15:414), lemma_15_5_holds (Proofs15Restriction:3116) (exact)	degree-weighted Fejér selection, field-size exponent $E_{k,m} \sim 2s/k$, retention $\alpha^s \eta^{s-1}/(2(s + 2)^s)$ (35)	not formalized; explicit N_0 for lemma_15_5 a marked question
Lemma 16.10	lemma_16_10 (Section16:851) open ; Section16SampledOrAnchoredOn (Proofs16ShortColumns:27)	sampler keeping sampled targets, $2q(m - r)/(m(r + 1))$; closed finite minimax for its rule; labelled refinement (39; the sampler is source 19's)	not formalized; candidate; the closing comparison stays open

Gowers	Lean statement; companion (status)	Claim of sources 34–39	Improvement
Props. 17.2, 17.7	<code>proposition_17_2</code> (Sections17_18:48), <code>proposition_17_2_holds</code> (Proofs17PhaseRemoval:373); <code>proposition_17_7</code> (Sections17_18:140), <code>proposition_17_7_holds</code> (Proofs17LocalizedPhaseRemoval:2941) (exact)	sharp symmetry threshold $1 - (1 - 1/p)^k$ over $\mathbb{F}_p^n$; lossless polarization for $p > k + 1$ (36)	not formalized; candidate (rank-two forms)

[write] **6 October 2026; sources 34–39.** *Formal status after the seventh write.* No statement of sources 34–39 is formalized, and none gains formal status by its placement; the ledger is unchanged. The only formalized statement of the report is still the lifting lemma of source 04. Two exact formal statements are now implied by written results of this report, which changes nothing formal because both already have exact companions: `lemma_5_2` by source 37’s Corollary I.55.2, and `lemma_5_5` (for $t \geq \text{weylThreshold } k$, with the formal Weyl input) by its Corollary I.57.3 (notes there). Formalization candidates, none claimed: source 37’s lower-bound certificate for Lemma 5.2 and its recurrence conversion; source 39’s sampler at the `Section16SampledOrAnchoredOn` interface (source 19’s first); source 36’s bilinear estimate for rank-two forms; source 34’s exact reduction, small-prime certificates and fourth-norm inequality (its own proposal). Sources 34–39 ship no Lean files; the `FORMALIZATION_PLAN.md` files of sources 36 and 39 and source 37’s `INTEGRATION.md` (shipped as `36-symmetry-certification-FORMALIZATION_PLAN.md`, `39-shared-anchor-FORMALIZATION_PLAN.md`, `37-missing-intervals-INTEGRATION.md`) are plans.

Re-check of the earlier tables at 89c1cbc38. Since the sixth write’s HEAD 5d4718772 the tracked Lean files of the development have changed by new modules, a facade, six lines of `Proofs15ProductToRestriction.lean` and thirty lines added to `Proofs15Restriction.lean` (two private declarations made public, one new theorem). The script of the write re-read the references of the six earlier tables at 89c1cbc38 and found the items recorded at the fifth and sixth writes (the hand-resolved list-pairing artifacts and the move of `proposition_17_7_holds` to line 2941) and two new moves: `lemma_15_5_holds` from line 3086 to line 3116 and `selectionDegenerate_real_small` from line 2690 to line 2692 (fourth and fifth tables); the dated tables keep their numbers.

[write] **7 October 2026; sources 40–46.** *The ledger at the eighth write.* At HEAD 6d4942222 (7 October 2026) the ledger records 113 exact companions and 7 open statements, and `FORMALIZATION_STATUS.txt` reports 713 modules and 4,583 audited public theorems; since the seventh write’s HEAD 89c1cbc38 `corollary_5_8` (b548d998e), `theorem_13_12`, `lemma_13_9`, `lemma_13_5`, `lemma_13_6` and `corollary_13_10` have closed. The open statements are `theorem_1_3`, `lemma_16_10`, `theorem_16_2`, `corollary_16_11`, `theorem_18_1`, `theorem_18_2` and `corollary_18_7`. At the pins of sources 41–45 it recorded 99 and 21 (c94998a5d, source 41) or 101 and 19 (b16ce380d, 797aa3cfe, 8c12c906e, 6ce33539e); sources 40 and 46 record no commit. Closed since those pins, besides the six above: `theorem_1_2`, `theorem_7_1`, `theorem_7_2`, `theorem_8_2`, `corollary_8_3` and `lemma_16_4` (and, since source 41’s pin, `lemma_13_7` and `lemma_13_8`); of these only `corollary_5_8` (sources 41 and 42) and `theorem_7_1`, `theorem_7_2` (source 45’s energy material, which compares with their statements, not their status) touch a subject of sources 40–46. No source of this write quotes ledger counts. Every Lean file and every paper file that sources 40–46 cite has the same blob at HEAD as at their pins (the intake’s provenance check, rerun at the write: 30 records, all matching).

Corollary 5.8 is closed. Since the seventh write the repository has formalized, in a sequence of commits from 4ac12a026 to b548d998e (6 October 2026, the later ones while this write was in progress), the two-regime architecture of sources 26, 32 and 41 with the catalogue’s constants: the cyclic progression second moment and the bound on nonproper directions, the variance increment `exists_proper_progression_density_increment` (`Proofs05VarianceIncrement.lean`, line 89; source 41’s Lemma I.66.1, hypothesis $L^3 \leq N$, its docstring crediting Report277), its low-correlation form, the cover-mass bounds, the weighted phase

transfer, the high-correlation assembly and parent scale, the target rounding $(4L)^{K_k} < N/M$, the reduction `corollary_5_8_holds_of_threshold_free_localization` to `ThresholdFreePolynomialLocalization` k , the affine and quadratic localizations with `corollary_5_8_degree_one` and `corollary_5_8_degree_two`, a finite recurrence in every degree, and `thresholdFreePolynomialLocalization_all`; the exact companion is `corollary_5_8_holds`. The localization uses the repository’s own finite Weyl estimate (radius exponent $2((k-1)^2+1)(2^{k-1}+3)$), not source 41’s recurrence ($B_k = 2^k + 8$); `FORMALIZATION_STATUS.txt` names source 41 as giving “a related stronger written localization budget”. The formal theorem has the catalogue’s constants (length $(N/M)^{1/K_k}/8$, gain $\alpha/16$); source 41’s stronger Theorem F.46.1 is not formalized. The rows below give file and line numbers at 6d4942222; every one was checked by reading the declaration at that line, and every comparison of sources 40–46 with a formal statement was checked against these declarations.

Gowers	Lean statement; companion (status)	Claim of sources 40–46	Improvement
Lemma 3.10	<code>lemma_3_10</code> (Sections01_03:340), <code>lemma_3_10_holds</code> (Proofs03CubeUpper:82) (exact)	$c_4^{\text{odd}} \leq (832/4795)^{1/4}$, quartic coefficient $P_d(832/4795)^{1/4}$ for $d \geq 4$ (45)	not formalized
Cor. 5.8	<code>corollary_5_8</code> (Section05:242), <code>corollary_5_8_holds</code> (Proofs05 ThresholdFreeLocalization:86) (exact, new); <code>corollary_5_8_holds_of_thres hold_free_localization</code> (Proofs05ThresholdFreeCorolla ry:145), <code>ThresholdFreePolynom ialLocalization</code> (Proofs05Thre sholdFreeCorollary:16), <code>thresholdFreePolynomialLocal ization_all</code> (Proofs05Threshol dFreeLocalization:70); <code>exists_proper_progression_de nsity_increment</code> (Proofs05Vari anceIncrement:89); <code>density_ce ll_of_high_correlation</code> (Proofs05HighCorrelationAssem bly:53); <code>corollary_5_8_degree_one</code> (Pr oofs05AffineLocalization:62); <code>corollary_5_8_degree_two</code> (Proofs05QuadraticLocalizatio n:87)	gain $\alpha/3$ at length $\max\{1, (N/M)^{1/D_k}/8\}$ in every degree, no scale premise (41); variance gain under $N \geq 2(L-1)^2 + 1$, affine cap $\sqrt{N/2}$ (42)	implies the formal conclusion and supplies its localization hypothesis on paper ([write] notes); the formal proof uses the same architecture and its own recurrence; 41’s constants and 42 not formalized
Section 7 (energy)	<code>theorem_7_1</code> (Sections06_07:104), <code>theorem_7_1_holds</code> (Proofs07FreimanClosure:185); <code>theorem_7_2</code> (Sections06_07:125), <code>theorem_7_2_holds</code> (Proofs07FreimanClosure:235) (exact)	three largest torsion-free energies; coset envelope for $\varepsilon < 1/9$ with all equality cases (45)	different statements; not formalized

Gowers	Lean statement; companion (status)	Claim of sources 40–46	Improvement
Lemma 15.4	<code>lemma_15_4</code> (Sections14_15:402), <code>lemma_15_4_holds</code> (Proofs15De <code>generateArrangements</code> :1131) (exact)	core directions for every $p \geq 5$; exact core products; $b_2(d)$ for every d , $b_2(8) = 3,280,234,631,541$ (44)	not formalized
Lemma 15.5	<code>lemma_15_5</code> (Sections14_15:414), <code>lemma_15_5_holds</code> (Proofs15Restriction:3116) (exact)	one Fejér kernel, retention $\alpha^s \eta^{s-1} / (3^{s-1}(s+1)^s)$, kernel-class barrier (40)	not formalized; explicit N_0 a marked question
Lemma 17.1, Prop. 17.2	<code>lemma_17_1</code> (Sections17_18:38), <code>lemma_17_1_holds</code> (Proofs17Phase:374); <code>proposition_17_2</code> (Sections17_18:48), <code>proposition_17_2_holds</code> (Proofs17PhaseRemoval:373) (exact)	over $\mathbb{F}_2^n$: exact cubic norm $2^{-r/2}$, codimension $r/2$, Tidor’s $8 \log_2(1/\delta)$ to $2 \log_2(1/\delta)$ (46, 43); certification $1 - 2^{1-d}$ and finite-abelian orthogonality (43)	outside the formal setting (N prime, $k+1 < N$); not formalized

[write] **7 October 2026; sources 40–46.** *Formal status after the eighth write.* No statement of sources 40–46 is formalized as stated, and none gains formal status by its placement; the ledger changes are the repository’s own work. The only formalized statement of the report itself is still the lifting lemma of source 04. The repository’s new modules prove `corollary_5_8` with the architecture of sources 26, 32 and 41 and the catalogue’s constants, not source 41’s theorem; the variance lemma they formalize is the one source 41 shares with sources 26 and 32, and their localization is the repository’s own (source 41’s is a second, written route). Formalization candidates, none claimed: source 41’s stronger theorem (exponent $1/D_k$, gain $\alpha/3$) along its recurrence and budgets; source 42’s arithmetic increment in place of $L^3 \leq N$; source 44’s width lemma and signed-collision product; source 45’s ordered-completion identity and rational U^4 certificate; source 43’s orthogonality theorem. The `FORMALIZATION.md` files of sources 43, 44 and 46 (shipped as `43-boolean-phase-FORMALIZATION.md`, `44-grid-core-FORMALIZATION.md`, `46-polarization-barriers-FORMALIZATION.md`) are plans; sources 40–46 ship no Lean files.

Re-check of the earlier tables at 6d4942222. Since the seventh write’s HEAD 89c1cbc38 the tracked Lean files of the development have gained 263 new files, and 24 existing files have changed (among them `Sections12_13.lean`, where the Section 13 statements were revised and closed). The script of the write re-read the 449 references of the seven earlier tables at 6d4942222 and found the items recorded at the fifth to seventh writes (the hand-resolved list-pairing artifacts and the recorded moves) and one new move: `lemma_13_5` from line 410 to line 420 of `Sections12_13.lean` (fourth table, two entries; the statement is now exact, closed in bc87bdf14); the dated tables keep their numbers and their dated statuses.

[write] **7 October 2026; sources 47–53.** *The ledger at the ninth write.* At HEAD 06b30ca52 (7 October 2026) the ledger records 113 exact companions and 7 open statements, unchanged since the eighth write’s HEAD 6d4942222, and `FORMALIZATION_STATUS.txt` reports 746 modules and 4,742 audited public theorems (713 and 4,583 then). The open statements are still `theorem_1_3`, `lemma_16_10`, `theorem_16_2`, `corollary_16_11`, `theorem_18_1`, `theorem_18_2` and `corollary_18_7`. At the pins of sources 47–53 (6 October 2026, 17:11–18:11) it recorded 101 and 19; the twelve statements closed since then are listed in Section F.5. None of sources 47–53 quotes ledger counts. Every Lean file the sources cite has its pin blob at HEAD except six, listed in the provenance (Section F.4) and annotated where they are cited.

Formal work on the subjects of sources 47–53 since their pins. *Lemma 16.10:* the finite-alphabet cover bounds, balanced words on long progressions, the identity-partition certificates of both packaged premises and the literal negation `lemma_16_10_packaged_premises_counterexample` (commits d6caf29c5 to c0a404a59, 7 October 2026, 00:21–00:47), along source 50’s route and crediting it; a positive explicit cover for

the packaged premises (1a8581c65); the endpoint equivalence of the full-domain unit product property with multilinearity (ff069f96e), source 50's Proposition II.234.1; the local affine lift rebuilt with a linear sampling budget and compressed candidates (93e36c6df, 081abc4d3); a conditional derivation of `corollary_16_11` from the Section 16 induction. *Section 13*: Lemmas 13.5, 13.6 and 13.9 and Corollary 13.10 closed with construction geometry carried forward (bc87bdf14 to 8da5b7f07), the older interfaces preserved; the square exponent improved at each fixed exact density (986165438); an explicit density-uniform threshold for square extraction (93d436414); Theorem 13.12 closed earlier (d48e8af00). The rows below give file and line numbers at 06b30ca52; every one was checked by reading the declaration at that line.

Gowers	Lean statement; companion (status)	Claim of sources 47–53	Improvement
Lemma 3.10	<code>lemma_3_10</code> (Sections01_03:340), <code>lemma_3_10_holds</code> (Proofs03CubeUpper:82) (exact)	$c_4^{\text{odd}} \leq (4/29)^{1/4}$, $c_5^{\text{odd}} \leq 29/50$ (49)	not formalized
Lemma 9.3	<code>lemma_9_3</code> (Sections08_09:77), <code>lemma_9_3_holds</code> (Proofs09Restriction:376) (exact)	constructive restriction by exact inventories, source 17's retention (48)	not formalized
Stage 13.7	<code>lemma_13_7</code> (Sections12_13:558), <code>lemma_13_7_holds</code> (Proofs13A1 1ScalesStepSpan:118) (exact); <code>lemma_13_7_without_domain</code> (Sections12_13:526), <code>lemma_13_7 _without_prime_assumption</code> (Sections12_13:534) (preserved, unproved)	the two preserved encodings are false over $N = 2^k$ (52)	negations not formalized; <code>lemma_13_7</code> untouched
Cor. 13.10	<code>corollary_13_10</code> (Sections12_13:686), <code>corollary_13_10_holds</code> (Proofs13ConstructedChain:31) (exact); <code>corollary_13_10_with out_construction_geometry</code> (Sections12_13:669) (unasserted); <code>Corollary1310Grid</code> (Proofs13Bi linearRestriction:33), <code>corollary_13_10_of_stage_cha in_and_grid</code> (Proofs13Bilinear Restriction:270)	retained data do not imply the grid; repair by reselection, exponent 2γ (51)	concerns the older interface; not formalized
Square extraction (Thm. 13.12)	<code>section13_complete_square_ex traction_uniform_density</code> (Proofs13UniformSquareExtract ion:15); <code>section13_complete_s quare_extraction_explicit</code> (Proofs13ExplicitGeometricThr eshold:24); <code>section13_complete_square_ex traction_below_limit</code> (Proofs 13SquareExponentLimit:63); <code>theorem_13_12</code> (Sections12_13:718), <code>theorem_13_12_holds</code> (Proofs13 EndpointPrintedTheorem:73) (exact)	closed threshold $N_0(a_0)$, side $D(a)N^{4\gamma(a)} - 1 \geq N^{3\gamma(a)}$, actual-density mass (53)	the repository's explicit threshold is another formula, at the lower density; 53 not formalized

Gowers	Lean statement; companion (status)	Claim of sources 47–53	Improvement
Lemma 15.4	<code>lemma_15_4</code> (Sections14_15:402), <code>lemma_15_4_holds</code> (Proofs15GenerateArrangements:1131) (exact)	two-cube $\delta_{k,1}$: second coefficient, exact for $k \leq 3$ (49)	another convention; not formalized
Lemma 16.10	<code>lemma_16_10</code> (Section16:851) (open); <code>lemma_16_10_packaged_premises_counterexample</code> (Proofs16PackagedPremisesCounterexample:13); <code>Section16AllBoxLineCovers.explicit_multilinear_cover</code> (Proofs16PackagedExplicitCover:11); <code>section16_unit_product_iff_multilinear</code> (Proofs16MultilinearProduct:89)	the catalogue statement is false; branch-cover laws; endpoint characterization (50)	negation and endpoint formalized (repository’s own work); the repair open
Cor. 16.11	<code>corollary_16_11</code> (Section16:893) (open); <code>corollary_16_11_of_theorem_16_2</code> (Proofs16CorollaryFromInduction:123)	acyclic replacement from I1–I7 (49)	conditional; not formalized
Prop. 17.2	<code>proposition_17_2</code> (Sections17_18:48), <code>proposition_17_2_holds</code> (Proofs17PhaseRemoval:373) (exact)	over $\mathbb{F}_2^n$: $\mathcal{M}(C_d) = 1 - (d+1)/2^d$, thresholds through degree six (47)	outside the formal setting; not formalized

[write] **7 October 2026; sources 47–53.** *Formal status after the ninth write.* No statement of sources 47–53 is formalized as stated, and none gains formal status by its placement; the ledger changes are the repository’s own work. The repository’s formal negation of `lemma_16_10` and its endpoint theorem follow source 50 and are credited to it in `FORMALIZATION_STATUS.txt`; they are the repository’s proofs, not this report’s. The only formalized statement of the report itself is still the lifting lemma of source 04. Formalization candidates, none claimed: source 52’s two negations; source 53’s closed threshold and actual-density exponent, with the exported length $\ell R \leq m + 1$; source 49’s translation reduction, rank-two classification and finite certificates; source 47’s even-period lemma and canonical induction; source 48’s inventory invariants. The `FORMALIZATION.md` files of sources 47 and 48 and source 50’s `FORMALIZATION_PLAN.md` (shipped with the prefixes `47-boolean-energies-`, `48-derandomization-`, `50-branch-cover-`) are plans; sources 47–53 ship no Lean files that are staged here.

Re-check of the earlier tables at 06b30ca52. Since the eighth write’s HEAD `6d4942222` the Lean development has gained 33 new files, and 7 existing files have changed, all in Section 16 (among them `Proofs16LocalAffineLift.lean`). The 495 references of the eight earlier tables were re-read at HEAD; the 23 flagged are exactly those recorded at the eighth write (the hand-resolved list-pairing artifacts and the recorded moves, among them `lemma_13_5`), and there is no new move. The dated tables keep their numbers and their dated statuses.

[write] **7 October 2026; sources 54–59.** *The ledger at the tenth write.* At HEAD `fe7165a3c` (7 October 2026) the ledger records 113 exact companions and 7 open statements, unchanged since the ninth write’s HEAD `06b30ca52`, and `FORMALIZATION_STATUS.txt` reports 755 modules and 4,836 audited public theorems (746 and 4,742 then). The open statements are still `theorem_1_3`, `lemma_16_10`, `theorem_16_2`, `corollary_16_11`, `theorem_18_1`, `theorem_18_2` and `corollary_18_7`. At the pins of sources 54–59 (6 October 2026, 19:52–20:29) it recorded 103 and 17 (source 54) and 105 and 15 (the others); `theorem_13_12` was open at

all of them. None of sources 54–59 quotes ledger counts. Every Lean file the sources cite has its pin blob at HEAD except `Sections12_13.lean` (source 57), whose cited definitions are unchanged and have moved (`theorem_13_12` from line 626 to line 718; provenance, Section F.4).

Formal work on the subjects of sources 54–59 since their pins. *Theorem 13.12* (source 57): the uniform quantitative Fejér purification and retained density (`e40138bd5`, 6 October 2026, 23:09), the propagation through the prepared frequency graph and the printed range away from one (`3e65784f4`, 23:18), source 57’s near-maximal endpoint and the exact companion `theorem_13_12_holds` (`d48e8af00`, 23:55), and the strengthening to length exponent 2^{61} with density exponent 2^{42} (`b3cb2dfa9`, 23:57); `FORMALIZATION_STAT_US.txt` names source 57’s route (`c6b9e8da9`, 22:06, before these commits). The explicit threshold of the uniform square extraction (`93d436414`) also bears on source 57’s threshold question. Nothing else on the subjects of sources 54, 55, 56, 58 and 59 has changed: `proposition_6_1`, `lemma_7_5`, `corollary_7_6`, `proposition_8_1`, `lemma_9_3`, `corollary_9_4` and `lemma_10_9` were exact at the pins and are unchanged.

Gowers	Lean statement; companion (status)	Claim of sources 54–59	Improvement
Prop. 6.1, Lemma 6.2	<code>proposition_6_1</code> (<code>Sections06_07:37</code>), <code>proposition_6_1_holds</code> (<code>Proofs05_10:505</code>); <code>lemma_6_2</code> (<code>Sections06_07:47</code>), <code>lemma_6_2_holds</code> (<code>Proofs05_10:589</code>) (exact)	sharp endpoint stability in every U^d (55); near-maximal U^k extraction (57); affine radii $4/9$, $1/2$, recovery below $2/9$, $\alpha^4 > 7/9$ (59); relative affine repair (56)	not formalized, except source 57’s U^4 endpoint in the repository’s own form (<code>endpoint_global_fourier</code> , <code>Proofs13EndpointGlobalFourier:12</code>)
Lemma 7.5, Cor. 7.6	<code>lemma_7_5</code> (<code>Sections06_07:183</code>), <code>lemma_7_5_holds</code> (<code>Proofs07FreimanRestriction:267</code>); <code>corollary_7_6</code> (<code>Sections06_07:197</code>), <code>corollary_7_6_holds</code> (<code>Proofs07AdditiveRestriction:100</code>) (exact)	exact separation ranks; restrictions over arbitrary abelian targets, polynomial in the doubling (59)	another setting (all abelian targets); not formalized
Prop. 8.1	<code>proposition_8_1</code> (<code>Sections08_09:29</code>), <code>proposition_8_1_holds</code> (<code>Proofs08QuadraticBias:1069</code>) (exact)	one common quadratic phase, coefficient $\geq 27/4$; offset ceilings and stability (54)	implies the catalogue statement on paper; not formalized
Lemma 9.3, Cor. 9.4	<code>lemma_9_3</code> (<code>Sections08_09:77</code>), <code>corollary_9_4</code> (<code>Sections08_09:91</code>); <code>lemma_9_3_holds</code> , <code>corollary_9_4_holds</code> (<code>Proofs09Restriction:376</code> , 416) (exact)	every cyclic modulus, exponent 54, explicit N_0 ; the positive-filter barrier (58)	implies the catalogue statements on paper; not formalized
Lemma 10.9	<code>lemma_10_9</code> (<code>Section10:335</code>), <code>lemma_10_9_holds</code> (<code>Proofs10Shift:717</code>) (exact)	an additional high-agreement branch on uniform domains, not a stronger Lemma 10.9 (56)	not formalized

Gowers	Lean statement; companion (status)	Claim of sources 54–59	Improvement
Thm. 13.12	<code>theorem_13_12</code> (Sections12_13:718), <code>theorem_13_12_holds</code> (Proofs13EndpointPrintedTheorem:73) (exact since d48e8af00); <code>theorem_13_12_away_from_one</code> (Proofs13FejerPrintedRange:39), <code>theorem_13_12_near_one</code> (Proofs13EndpointPrintedTheorem:12), <code>theorem_13_12_with_fejer_explicit_bound</code> (Proofs13FejerExplicitExponent:55), <code>theorem_13_12_with_improved_printed_bounds</code> (Proofs13ImprovedPrintedBounds:11)	the printed length with density $(\alpha/2)^{2^{41}}$ (57)	proved formally along this route with density $(\alpha/2)^{2^{42}}$; the 2^{41} not formalized
Section 13 threshold	<code>section13_complete_square_extraction_uniform_density</code> (Proofs13UniformSquareExtraction:15); <code>section13_complete_square_extraction_explicit</code> (Proofs13ExplicitGeometricThreshold:24)	N_{sq} left inexplicit (57), made closed by source 53's threshold	the repository's explicit formula is another one

[write] 7 October 2026; sources 54–59. *Formal status after the tenth write.* No statement of sources 54–59 is formalized as stated, and none gains formal status by its placement; the ledger changes are the repository's own work. The formal proof of `theorem_13_12` follows source 57's route and credits it in `FORMALIZATION_STATUS.txt`; it is the repository's proof, with its own constants, not this report's. The only formalized statement of the report itself is still the lifting lemma of source 04. Formalization candidates, none claimed: source 58's Theorem A as a strengthened `lemma_9_3`; source 54's common-phase theorem beside `proposition_8_1`; source 55's explicit local theorem; source 59's separation ranks and signed affine identities; source 57's density exponent 2^{41} and partition classification; source 56's relative repair. The `FORMALIZATION.md` and `FORMALIZATION_PLAN.md` files of sources 54 and 55 are plans; no source ships Lean code.

Re-check of the earlier tables at fe7165a3c. Since the ninth write's HEAD 06b30ca52 the Lean development has gained 9 new files, all in Section 16, and no existing file of the development has changed; so every file and line cited by the nine earlier tables reads as at the ninth write, where their 551 references gave the 23 recorded items and no new move. The dated tables keep their numbers and their dated statuses.

[write] 7 October 2026; sources 60–69. *The ledger at the eleventh write.* At HEAD fb9602e55 (7 October 2026) the ledger records 114 exact companions and 6 open statements (113 and 7 at the tenth write's HEAD fe7165a3c), and `FORMALIZATION_STATUS.txt` reports 764 modules and 4,877 audited public theorems in its last complete audit (755 and 4,836 at the tenth write). The change is the repair of Lemma 16.10 (456b6e3ad and 851eea533, 7 October 2026, 12:20 and 12:25): `lemma_16_10` now states the cover the printed proof establishes, and its companion `lemma_16_10_holds` (Proofs16Lemma10.lean, line 12) is counted but, as `FORMALIZATION_STATUS.txt` states, not yet kernel-checked (the complete audit above predates it and counted 113/7; for the repair only `Section16.lean` and 61 upstream modules were compiled). This write ran no Lean build and does not confirm it. The open statements are `theorem_1_3`, `theorem_16_2`, `corollary_16_11`, `theorem_18_1`, `theorem_18_2` and `corollary_18_7`. At the pins of sources 60–69 (6 October 2026, 21:05–22:07) it recorded 107 and 13; of the seven statements counted as exact since, the six closed before the tenth write's HEAD (`corollary_5_8`, `lemma_13_5`, `lemma_13_6`, `lemma_13_9`, `corollary_13_10`, `theorem_13_12`)

are not on the subjects of these sources, and the seventh, the repaired Lemma 16.10, bears on source 67's (below). None of sources 60–69 quotes ledger counts. Every Lean file the sources cite has its pin blob at HEAD (Sections06_07.lean, Sections08_09.lean, Proofs09Restriction.lean, Definitions.lean, Section04.lean, Sections17_18.lean; provenance, Section F.4).

Formal work on the subjects of sources 60–69 since their pins. Lemma 16.10 (source 67): the affine interpolation count, formalized as Section16SampledOrAnchoredOn.compressed_multilinear_cover (Proofs16CompressedInterpolation.lean, line 56) with its budget comparisons and the propagation through the Section 16 affine lift (081abc4d3, 7 October 2026, 08:49, nineteen minutes after source 67's arrival); FORMALIZATION_STATUS.txt names source 67's theorem as the origin and says that it does not prove the unit-parameter conclusion. Later the same day the catalogue entry was repaired (above): the printed encoding, refuted along source 50's route, is kept as lemma_16_10_printed_unit_encoding (Section16.lean, line 852), and the repaired lemma_16_10 (line 912) bounds the graph count by $\max(b, \binom{r}{2}b^2)$, the form of the compressed count, through Section16AllBoxLineCovers.explicit_multilinear_cover (Proofs16PackagedExplicitCover.lean, line 11). Nothing else on the subjects of sources 60–66, 68 and 69 has changed: lemma_5_2, proposition_6_1, lemma_6_2, lemma_7_5, corollary_7_6, lemma_9_3, corollary_9_4, lemma_12_3, lemma_3_10, proposition_17_2 and proposition_17_7 were exact at the pins and are unchanged; conjecture_4_1 and conjecture_4_2 are definitions, not asserted.

Gowers	Lean statement; companion (status)	Claim of sources 60–69	Improvement
Lemma 5.2	lemma_5_2 (Section05:141), lemma_5_2_holds (Proofs05MissingInterval:220) (exact)	unequal-budget missing-arc mass near the equal-budget branch (67)	the extremal problem behind it; not formalized
Prop. 6.1, Lemma 6.2	proposition_6_1 (Sections06_07:37), proposition_6_1_holds (Proofs05_10:505); lemma_6_2 (Sections06_07:47), lemma_6_2_holds (Proofs05_10:589) (exact)	affine rigidity into elementary abelian 3-groups on $t \leq 2/3$; recovery thresholds 1/4, 4/9; cutoff 13 (67)	not formalized
Thm. 7.2 (energy)	theorem_7_2 (Sections06_07:125), theorem_7_2_holds (Proofs07FreimanClosure:235) (exact)	energy gap with coefficient one, coset range 1/4, uniqueness 2/9 (68); partition normal forms, the fourth level (66)	no declaration concerns energy levels; not formalized
Lemma 7.5, Cor. 7.6	lemma_7_5 (Sections06_07:183), lemma_7_5_holds (Proofs07FreimanRestriction:267); corollary_7_6 (Sections06_07:197), corollary_7_6_holds (Proofs07AdditiveRestriction:100) (exact)	exact Freiman partitions on arbitrary finite abelian groups (61)	another setting (all finite abelian groups, prime moduli in the catalogue); not formalized
Lemma 9.3, Cor. 9.4	lemma_9_3 (Sections08_09:77), corollary_9_4 (Sections08_09:91); lemma_9_3_holds, corollary_9_4_holds (Proofs09Restriction:376, 416) (exact)	exponent 53 for arbitrary targets (68); $\alpha^{31}\eta^{30}$ for cyclic targets (62); threshold-free exact restriction (61); prime moduli at threshold L^{m-1} (60)	61, 62, 68 imply the catalogue statements on paper (61 with $N_0 = 1$, also corollary_9_4); not formalized

Gowers	Lean statement; companion (status)	Claim of sources 60–69	Improvement
Lemmas 12.3–12.6	<code>lemma_12_3</code> (Sections12_13:151), <code>lemma_12_3_holds</code> (Proofs12HigherArrangements:334); <code>lemma_12_5</code> , <code>lemma_12_6</code> (Sections12_13:170, 183; exact)	the arrangement polynomial of every order, equality cases, two-direction repair (60)	full products only; not formalized
Lemma 16.10	<code>lemma_16_10</code> (Section16:912) (repaired; <code>lemma_16_10_holds</code> , <code>Proofs16Lemma10:12</code> , not yet kernel-checked); printed encoding <code>lemma_16_10_printed_unit_encoding</code> (Section16:852) refuted	interpolation count $E_t(L_1, \dots, L_r)$, exact for the interface (67)	affine case formalized: <code>Section16SampledOrAnchoredOn.compressed_multilinear_cover</code> (<code>Proofs16CompressedInterpolation:56</code>)
Lemma 3.10	<code>lemma_3_10</code> (Sections01_03:340), <code>lemma_3_10_holds</code> (<code>Proofs03CubeUpper:82</code>) (exact)	$\lim c_d^{\text{odd}} = 1/3$ with a rate; the quartic coefficient $P_d c_d^{\text{odd}}$ (63; 60)	not formalized
Conj. 4.1, 4.2	<code>conjecture_4_1</code> , <code>conjecture_4_2</code> (Section04:30, 39) (definitions, not asserted)	binomial separation, $\rho_4 \ll \alpha^5$, $\rho_6 \ll \alpha^{7.847}$ (64)	no formal counterpart
Prop. 17.2	<code>proposition_17_2</code> (Sections17_18:48), <code>proposition_17_2_holds</code> (<code>Proofs17PhaseRemoval:373</code>); <code>lemma_17_1</code> (Sections17_18:38), <code>lemma_17_1_holds</code> (<code>Proofs17Phase:374</code>) (exact)	the Boolean threshold $1 - (d+1)/2^d$ in every degree (65, 69)	characteristic two, outside the prime-cyclic statements; not formalized
Prop. 17.7	<code>proposition_17_7</code> (Sections17_18:140), <code>proposition_17_7_holds</code> (<code>Proofs17LocalizedPhaseRemoval:2941</code>) (exact)	local coefficient asymptotic to $2^{k+1}/(k+2)!$ (63)	not formalized

[write] **7 October 2026; sources 60–69.** *Formal status after the eleventh write.* No statement of sources 60–69 is formalized as stated, and none gains formal status by its placement; the one ledger change since the tenth write is the repair of Lemma 16.10 (above), and the formalizations on their subjects, the affine interpolation count of source 67 and the graph count of the repaired lemma, are the repository’s own, in its own form. The only formalized statement of the report itself is still the lifting lemma of source 04. Formalization candidates, none claimed: source 68’s exponent-53 theorem and source 61’s threshold-free exact restriction as strengthened `lemma_9_3`; source 68’s energy gap and coset theorem; source 66’s partition normal form; source 65’s integration criterion and endpoint recognition; source 63’s bounded inequality and face-cover identity; source 64’s seed certificate; the higher-degree case of source 67’s count. The `FORMALIZATION.md` and `FORMALIZATION_PLAN.md` files of sources 61, 62, 64, 65 and 69 are plans; no source ships Lean code.

Re-check of the earlier tables at fb9602e55. Since the tenth write’s HEAD `fe7165a3c` the Lean development has gained ten files and seven existing files have changed, all in Sections 16 and 18. One of them is cited by the earlier tables: `Section16.lean`, changed by the repair of Lemma 16.10. Its lines 691, 749 and 824 (`Section16LineCover`, `lemma_16_9`, `Section16AllBoxLineCovers`) are unchanged; the predicate that the earlier tables and notes cite as `lemma_16_10` at line 851, open and refuted along source 50’s route, is now `lemma_16_10_printed_unit_encoding` at line 852, unchanged in content, and is the one meant wherever they say that `lemma_16_10` is open or false (its negation, `Proofs16PackagedPremisesCounterexample.lean`, moved from line 13 to line 14); `corollary_16_11` moved from line 893 to line 976 and is still open. Otherwise those tables read as at the tenth write and keep their numbers and dated statuses. Two further citations

elsewhere in the report moved (the ninth write’s note after Corollary II.235.11: `Section16ContextualLiftAt` from line 18 to line 20 of `Proofs16ContextualInduction.lean`, `theorem_16_2_of_selected_lifting` from line 64 to line 65 of `Proofs16SelectedLiftInduction.lean`), and the universal contextual-lift contract cited there has since been proved false (`section16_contextual_lift_counterexample`, `Proofs16ContextualLiftCounterexample.lean`, line 80; 48efbb992, after source 87, placed and not yet written in); Theorem 16.2 is not refuted, and the existential selected-lifting route stays open (dated note there).

V.2 Further questions and research [write]

Added in the write, 6 October 2026.

This section applies Vladimir’s standing rule of 4 October 2026: no unproved or wrong claim is dropped; open claims are stated as questions with their source, the argument they came with and what is missing; wrong claims stay on record with a proof of their failure. The five sources ask 46 questions: source 01 ten (Section V.4, Questions 1–10 there in order), source 02 seven (Section V.9, delivered Subsections 11.1–11.7), source 03 nine (Section V.14, Research questions 1–9 there in order), source 04 twelve (Section V.17, Q1–Q12) and source 05 eight (Section V.23, paragraphs 1–8). They are printed there as delivered. Below they are grouped by theme, with every asking source credited; the counts per theme are 14, 9, 8, 7 and 8. Items are numbered within their theme, so later additions extend a list without renumbering it.

[write] 6 October 2026; sources 06–10. *Sources 06–10.* They ask 50 further numbered questions: source 06 eleven (Section V.27, Questions 1–11 there), source 07 eight (Section V.29, Questions 1–8 there in order), source 08 ten (Section V.34, Questions 1–10 there in order), source 09 ten (Section V.39, Research questions 1–10, its Subsections 14.1–14.10) and source 10 eleven (Section V.44, items 1–11 there). Source 10 also ends its Sections III.29 and II.21 with a paragraph of questions; both are credited below. The intake dossier counted ten questions for source 10 and 95 in all; the count is 96. They are added to the themes below as new items or, where they ask the same thing, merged into an existing item by a dated sentence. The counts per theme become 17, 19, 33, 13 and 14 (the headings below keep the counts of sources 01–05).

Unproved or wrong claims in sources 06–10. The intake and the second write found no claim of sources 06–10 that is wrong, and no claim asserted as proved without a proof. One wrong statement is on record: the quartic lower bound with coefficient 12, which sources 08 and 09 pose and refute themselves and source 03 also refutes; Remark W1 keeps it with an explicit counterexample. Inputs used but not proved in the sources: Zuev’s threshold-function count (source 06, as for source 01), the Plünnecke–Ruzsa inequality (source 10, via Petridis) and, for source 10’s labelled modern-input variant only, the Reiher–Schoen theorem; these are published theorems. Source 10’s “independent review” is not shipped, and nothing rests on it. Statements of the intake corrected at the write: source 06’s improvement over the $M - 1$ corollary is strict only for $\alpha < 1$ (note after Corollary I.9.3); source 01’s increment is at least $3/2$ times source 10’s D_γ for every $\gamma \leq 1/2$ (exactly when $\gamma(4 - (2 - \delta)\kappa) \leq 2$), but not along the whole frontier (corrected after the independent check of 6 October 2026; the first wording read “source 01’s increment dominates source 10’s by $3/2$ only for $\gamma \leq 1/2$ ”, which is false as worded, since $2/(4 - (2 - \delta)\kappa) > 1/2$ for every $\kappa > 0$) (note after Table 3); source 09’s date for version 2 of Shi–Dong is wrong (bibliography).

Answered or advanced inside the report. Source 08’s question whether $\kappa_3^{\text{odd}} = 1/\sqrt{2}$ is answered yes by source 04 (item III.1). Source 03’s question on the fifth coefficient is answered by source 09 (item III.2). Source 02’s finite-cell problem is solved by source 06 in the weighted model with M fixed and the cell masses free; prescribed cell sizes and integer cardinalities remain open (item I.6) (corrected after the independent check of 6 October 2026; the first wording read “is solved in the weighted model by source 06”). Source 04’s Question 8 is answered a second time by source 06 (item IV.1). Source 06’s Question 7 is advanced by source 10 for one distinguished function (item III.9). Source 04’s Question 10 (non-admissible slopes) is advanced by source 10’s torsion factors (item III.7).

[write] 6 October 2026; sources 11–14. *Sources 11–14.* They ask 45 further questions: source 11 nine

(Section V.47, Research questions there in order, “11 Q1”–“11 Q9” below), source 12 twelve (Section V.51, its Questions in order, “12 Q1”–“12 Q12”), source 13 twelve (Section V.54, “13 Q1”–“13 Q12”) and source 14 twelve (Section V.57, the items of its research program, “14 Q1”–“14 Q12”). They are added to the themes below as new items or merged into existing items by dated sentences. The counts per theme become 24, 40, 44, 13 and 20 (141 questions; the headings keep the counts of sources 01–05). One further question is the intake’s, not a source’s, and is marked as such (item II.21).

Unproved or wrong claims in sources 11–14. The intake and the third write found no claim of sources 11–14 that is wrong and no claim asserted as proved without a proof. Recorded at the write: source 11’s cutoff 34, which it calls sufficient, cannot simply be dropped for general groups (Remark W2, an explicit counterexample to that tempting extension, not to a claim). Framing corrected in dated notes: source 13 presents its Fourier increment as a gain without noting that source 01’s is larger at the same length; source 12’s abstract says that its radius selection “replaces” Corollary 10.11, while the ambient radius is selected; source 12’s list of open obligations is incomplete; source 14 misplaces one declaration. Inputs used and not proved: the component-extraction lemmas (12, borrowed with stated hypotheses; formal), Warning’s second theorem and the Reed–Muller minimum-weight theorem (14 proves what it needs and credits them), the classical three-halves lemma (14, proved). Numerical evidence kept as evidence: source 14’s coefficient $2^{3/4}/3$ on groups of exponent five (item III.23). Source 14’s internal review is not shipped, and nothing rests on it.

Answered or advanced inside the report. Source 04’s Question 1 is settled by source 14 (Remark W4; item III.3). Source 04’s Question 2 is answered in U^2 , not in L^2 , by source 14 (item III.4), and its Question 7 is advanced by source 14’s arrangement equality (item II.3). Source 11’s Question 3 is answered in substance by composing it with source 14 (Remark W3; item II.20). Source 13’s Question 5 is solved by source 06 in the weighted model with M fixed and the cell masses free (item I.6). Source 05’s anisotropic question is advanced by source 13’s obstruction with independent tolerances (item I.8).

[write] 6 October 2026; sources 15–20. *Sources 15–20.* They ask 74 further questions: source 15 nine (Section V.60, its Questions in order, “15 Q1”–“15 Q9” below), source 16 fourteen (Section V.67, its items 1–14, “16 Q1”–“16 Q14”), source 17 ten (Section V.71, its paragraphs 1–10), source 18 twelve (Section V.78, its items 1–12), source 19 twelve (Section V.82, its Research questions in order, “19 Q1”–“19 Q12”) and five more at the end of its Section I.34 (“19 S6(1)”–“19 S6(5)”), and source 20 twelve (Section V.84, its Questions 1–12). They are added to the themes below as seventeen new items or merged into existing items by dated sentences. The counts per theme become 40, 68, 60, 17 and 30 (215 questions; the headings keep the counts of sources 01–05). Two further questions are the intake’s, not a source’s, and are marked as such (items II.6 and II.8).

Unproved or wrong claims in sources 15–20. The intake and the fourth write found no claim of sources 15–20 that is wrong and no claim asserted as proved without a proof. Corrected in dated notes, with the first wording kept as delivered: overstated novelty (source 15’s “sharp density refinements” and Cartesian analysis are source 04’s (A5)–(A11) and are superseded by source 14’s sharp constant; source 16’s Section 5 repeats source 04’s Lemmas 8.5–8.6 without credit, and its cube coefficients, like source 18’s Section 3, repeat source 09’s theorem; source 20’s “not another reordering” overlooks source 06’s quadratic estimate, which it could not see); understated formal status (sources 15, 16, 17: notes in Sections F.20, F.21 and F.22); stale repository claims (sources 16, 19); and two over-broad statements of source 17 (Section II.64: the multiplicity product is a valid lower bound, and the composite-modulus failure needs a small prime factor). Inputs used and not proved: Maynard’s published recurrence theorem and Lau’s preprint (sources 16, 18 and 19; the polynomial-in- q exponents of sources 18 and 19 and the own-step exponent of source 16 are conditional on them), and the classical cube circuits (source 18 credits Beerenwinkel–Pachter–Sturmfels). Source 19’s reading of printed page 575 of the 2001 paper was not checked at the write and is kept as its reading. The internal reviews mentioned by sources 18 and 19 are not shipped, and nothing rests on them.

Answered or advanced inside the report. Source 13’s Question 3 (item I.1) is answered by sources 18 and 19, conditionally on Maynard’s theorem. Source 06’s Question 7 is answered for the norm-conversion step by source 20 (item III.9), and its Question 6 is advanced by source 20’s strictly smaller coefficient (item III.8). Source 15’s first question is answered in its leading constant by source 14 (item II.31), and its second and third are source 14’s Questions 5–6. Source 17’s Question 7 is partly answered by source 10’s Corollary II.21.2. Source 19’s Question 8 is partly answered by source 18 (squaring is paid once, at the base;

item I.12). Source 18's Question 7 is settled at $d = 3$ by source 04's norm gap (item III.1).

Still pending. Sources 21–23 (placed by a902613c9), 24–27 (79d9075b9) and 28–33 (c5513046e; the last two placed while this write was in progress) are in this directory and not used in this text. Their placements record, among other findings: source 21 as answering source 14's exponent-five questions (item III.23); source 22 as treating source 18's $d = 1$ case, sources 25 and 27 as re-deriving source 18's leading coefficient and source 27 as answering source 18's first question for $d = 2$ (item II.38); source 23 as answering source 13's Question 1 (item I.12); sources 28 and 29 as bounding c_d^{odd} for $d \geq 4$ (item III.1) and source 29 as adding a fixed-field restriction theorem that keeps every degenerate tuple (item II.6). Those findings are for the writes that add them, and the items below still state the questions as open.

[write] 6 October 2026; sources 21–27. *Sources 21–27.* They ask 74 further questions: source 21 eleven (Section V.88, its Questions in order, “21 Q1”–“21 Q11” below), source 22 fifteen (Section V.92, its Research questions in order), source 23 twelve (Section V.94, its items 1–12), source 24 ten (Section V.98, its Questions in order), source 25 eight (Section V.103, its Questions 1–8), source 26 six (the items at the end of Section V.105) and source 27 twelve (Section V.108, its Research questions in order). They are added to the themes below as eight new items or merged into existing items by dated sentences. The counts per theme become 49, 104, 79, 17 and 40 (289 questions; the headings keep the counts of sources 01–05).

Unproved or wrong claims in sources 21–27. The intake and the fifth write found no claim of sources 21–27 that is wrong and no claim asserted as proved without a proof. Corrected in dated notes, with the first wording kept as delivered: overstated novelty or framing (source 22's “direct Section 15 improvement” gains nothing at Lemma 15.5 at leading order; source 24's table presents as its contributions results of sources 03 and 10, some with weaker constants, and its Proposition 7.3 “in every abelian group” is already formal; sources 25 and 27 present source 18's leading coefficient and single-certificate bound as their own, and source 27's remark on the research directory overlooks sources 18 and 22; all three could not see those texts), and the intake's description of source 23's sampling lemma as “a weaker version” of source 19's, made precise (its stated sample size is one larger; its displayed inequality bounds a different loss). Inputs used and not proved in the sources: the Reiher–Schoen theorem (source 24's Corollary II.80.2, marked as imported), and the classical results that source 22 credits (Warning's second theorem, the minimum-weight theorem of Delsarte, Goethals and MacWilliams). Source 27's exact eight-vertex count rests on a finite rational certificate, reproducible by its verifier and confirmed independently at intake, not checked by a proof assistant. Kept as evidence only: the intake's numerical convergence to 20/9 on $\mathbb{Z}/5\mathbb{Z}$, its exact values of $\delta_{k,1}(p)$ for $k \leq 3$, and its Monte Carlo at $k = 2$, $d = 2$. The reviews mentioned by sources 21 and 22 are not part of the argument here.

Answered or advanced inside the report. Source 14's Question 1 is answered by sources 21 and 23 (item III.23). Source 13's Question 1 is answered by source 23 (item I.12). Source 18's Question 1 is answered for $d = 2$ by source 27, and with it the same question of sources 22 and 25 (their first questions); source 18's Question 5 is answered for $d = 1$ in its leading coefficient by source 22, and the first-order count in small characteristic, source 25's Question 4, by source 27 (item II.38). Source 27's question on the single-pair regime is answered in its leading coefficient by source 22 (same item). Source 19's Question 6 is advanced by source 23 (item II.40). Source 14's Question 7 is advanced and its Question 8 has a first instance by source 22 (items III.19, III.20). Source 04's Question 6 and source 10's Question 6 are advanced for codomains $\mathbb{F}_p^d$ by source 24 (items II.8, II.17).

Still pending. Sources 28–33 (placed by c5513046e), 34–39 (placed by bf2fe146d) and 40–46 (placed by 2dfe46f44; both batches while this write was in progress) are in this directory and not used in this text. Their placements record, among other findings: sources 28, 29 and 34 as bounding c_d^{odd} for $d \geq 4$ (item III.1); source 29 as adding a fixed-field restriction theorem that keeps every degenerate tuple (item II.6); source 32 as proving Corollary 5.8 at degree two without a scale threshold (item I.19); source 33 as continuing the arrangement material of sources 14 and 15; source 34 as answering source 07's Question 1 (item II.9) and source 09's question on the prime-onset function (item III.16); source 35 as answering source 23's question on polynomial phases (item I.15); source 38 as reducing the exponent-five problem to $\mathbb{Z}/5\mathbb{Z}$ at small u/δ , answering three of source 21's questions (item III.27); source 39 as re-deriving source 19's sampling minimax (item II.39); source 40 as replacing the repeated Riesz products of Lemma 15.5 by one Fejér kernel (item II.44); source 41 as proving the whole of Corollary 5.8 in every degree with no scale premise, and source 42 as replacing the cube-root length cap by a square root (both item I.19); sources 43 and 46 as continuing source 36

in characteristic two; source 44 as computing the second Whitney coefficient of the cross-section arrangement in every dimension, answering two of source 27's questions (item II.38); and source 45 as improving the upper bound for c_4^{odd} (item III.1) and extending the range of source 14's coset envelope (item II.31). Those findings are for the writes that add them, and the items below still state the questions as open.

[write] 6 October 2026; sources 28–33. *Sources 28–33.* They ask 69 further questions: source 28 fifteen (Section V.113, its Research questions in order, “28 Q1”–“28 Q15” below), source 29 thirteen (Section V.117, its Questions in order), source 30 twelve (Section V.119), source 31 thirteen (Section V.123), source 32 four (the items at the end of Section V.126) and source 33 twelve (Section V.127). They are added to the themes below as four new items or merged into existing items by dated sentences. The counts per theme become 53, 150, 88, 17 and 50 (358 questions; the headings keep the counts of sources 01–05).

Unproved or wrong claims in sources 28–33. The intake and the sixth write found no claim of sources 28–33 that is wrong and no claim asserted as proved without a proof. Corrected or qualified in dated notes, with the first wording kept as delivered: source 28's pin of source 07 (39 digits, a 6 missing); source 29's reference to “Section 8” of source 17 (it is Section 7); source 28's “strictly improves the previously available $2^{-1/2}$ ” (true at its pin, superseded the same day by source 29); priority and credit: source 30 re-derives source 12's tree budget and compatibility inequality without citing it and has a weaker extension budget, source 31's Section 7 repeats source 10's mechanism with weaker constants and its Reiher–Schoen formula, and sources 28, 29 and 31 overlap one another from one pin without citations. Moved to questions as a marked question (not proved here): the intake's observation that source 28's theorem with its Corollary II.126.2 gives the formal **theorem_10_13** with better parameters (item II.47). Kept as evidence only: the intake's numerical optimum near 10.25714 for the U^4/U^2 ratio and its recurrence search for source 32's Lemma I.45.3. A review result, printed and credited to the intake: the certificate $c_4^{\text{odd}} \geq 0.5587725$ on $\mathbb{Z}/7\mathbb{Z}$, recomputed exactly at the write. Inputs used and not proved in the sources: Reiher–Schoen's Theorem 1.2 (source 31's Appendix A and modern comparison), Plünnecke–Ruzsa through Petridis (source 31), the classical results the sources credit.

Answered or advanced inside the report. Source 03's Research question 1, source 04's Question 3, source 08's question on κ_d^{odd} , source 14's Question 11, source 18's Question 7 and source 21's Question 7 are advanced for $d \geq 4$ by sources 28 and 29, and the sub-question whether real parts of characters are extremal is answered (no) for $d \geq 4$ by source 28; source 29's Question 1 (its second sentence) is answered by source 28 and source 28's question on a positive lower bound uniform in d by source 29 (item III.1). Source 02's Subsection 11.1 is answered in a prime-cyclic, bounded-fibre form by source 31 (item II.1) and its Subsection 11.2 advanced by sources 31 and 29 (item II.2). Source 04's Question 4 and source 17's Question 2 are answered for fixed fields by source 29 (item II.6). Source 12's Questions 4–6 are advanced by source 30 (item II.28). Source 11's Question 7, source 14's Questions 5–6 and source 15's Questions 2–3 are advanced on full products by source 33 (item II.24). Source 26's Question 4 is answered at degree two by source 32 (item I.19). Source 10's Question 7 is advanced by source 31's Appendix A (item II.18).

Still pending. Sources 34–39 (placed by **bf2fe146d**) and 40–46 (placed by **2dfe46f44**) are in this directory and not used in this text. Their placements record findings for items III.1 (sources 34, 35 and 45: upper coefficients 81/16, 5, 4795/832), I.19 (sources 41 and 42), II.46 (source 38: the odd-order stability modulus) and others listed at the fifth write; the items below still state the questions as open.

[write] 6 October 2026; sources 34–39. *Sources 34–39.* They ask 73 further questions: source 34 eleven (its Research questions in Section V.130, “34 Q1”–“34 Q11” below), source 35 fifteen (the Research question in Section III.72 as “35 Q1” and the fourteen of Section V.133 as “35 Q2”–“35 Q15”), source 36 twelve (Section V.136), source 37 ten (its Questions in Section V.139), source 38 fifteen (the numbered items of Section V.144) and source 39 ten (Section V.147). They are added to the themes below as six new items or merged into existing items by dated sentences. The counts per theme become 64, 169, 111, 28 and 59 (431 questions; the headings keep the counts of sources 01–05).

Unproved or wrong claims in sources 34–39. The intake and the seventh write found no claim of sources 34–39 that is wrong and no claim asserted as proved without a proof. Corrected or qualified in dated notes, with the first wording kept as delivered: priority and credit, namely source 34's reuse of source 28's Section 8 (two-harmonic example, witness, dual identity, single-pair non-extremality) without citation, which makes its “single-pair extremality left open” and “improves both endpoints” accurate only against source 29;

source 39’s re-derivation of source 19’s sampling theory without citation; source 36’s sharpening of source 31’s noncommuting-subgroup gap and its analogue of source 29’s symmetry step without citation; source 35’s “strongest inspected coefficient” (superseded the same day by source 34); source 35’s “strengthening” of the $1/(7\nu + 1)$ obstruction, which holds for prescribed lengths, not for source 19’s flexible-length obstruction; source 38’s ledger counts (those of its pin). A claim refined at the write: source 39’s closed solution of the finite profile problem holds for its own certification rule; for source 19’s rule the same formula fails in 199 of 2,835 small cases (exact counterexample $m = 2, r = 1, t = 2, q = 2$: formula $1/2$, true value 1; note after Theorem II.155.3), so source 19’s Question 1 stays open for its rule. Moved to questions as a marked question (not proved here): the intake’s observation that source 35’s Corollary II.145.3 gives the formal `lemma_15_5` with an explicit threshold (item II.50). Kept as evidence only: the intake’s numerical local minima of the U^4/U^2 ratio near 10.2571 ($c_4^{\text{odd}} \approx 0.55878$), its numerical confirmation of source 38’s u^{10} and u^{12} coefficients on $\mathbb{Z}/5\mathbb{Z}$, and the best member 0.5587462 of the two-harmonic family. Inputs used and not proved in the sources: the Alon–Vu unimodular matrices and Shitov’s composition (source 34, credited), the Weyl inequality in the formal form of `lemma_5_3` (source 37), classical positive-polynomial and truncated-moment theory (source 37, Yang–Xie credited), finite Heisenberg representation theory (source 36, proved in the needed form), Eisner–Tao’s classification (source 38, context).

Answered or advanced inside the report. Source 09’s Research question on the prime-onset function is answered by source 34 (item III.16); source 07’s Question 1 is answered in its finite-fibre model by source 34’s Appendix A (item II.9); item III.1 is advanced to $0.5587725 \leq c_4^{\text{odd}} \leq 2/3$ by source 34 (and source 35); source 21’s questions on the universal eighth-order coefficient, fixed higher dimensions and further terms on the five-point group are answered by source 38 (item III.27); source 31’s odd-order coefficients by source 38 (item II.46); source 23’s Question 6 in its construction half by source 35 (item I.15); source 19’s Question 1 for source 39’s rule and its rates question in part, by source 39 (item II.39); source 22’s question on defect formulas for more odd norm layers is advanced for three layers by source 38 (item III.28).

Still pending. Sources 40–46 (placed by `2dfe46f44`) are in this directory and not used in this text. Their placement records findings for items III.1 (source 45: $c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$), I.19 (source 41: a written proof of the whole Corollary 5.8; source 42: square-root scale), II.44 and II.50 (source 40: the one-stage Fejér selection that source 35 builds on; source 44: the second Whitney coefficient), the new Theme IV items (sources 43 and 46: characteristic two) and others; the items below still state the questions as open.

[write] 7 October 2026; sources 40–46. *Sources 40–46.* They ask 62 further questions: source 40 eight (its Research questions in Section V.152, “40 Q1”–“40 Q8”), source 41 five (the enumerated items of Section V.155), source 42 four (Section V.156), source 43 ten (Section V.160), source 44 nine (Section V.163), source 45 sixteen (Section V.166) and source 46 ten (Section V.171). They are added to the themes below as five new items or merged into existing items by dated sentences. The counts per theme become 72, 192, 115, 46 and 68 (493 questions; the headings keep the counts of sources 01–05).

Unproved or wrong claims in sources 40–46. The intake and the eighth write found no claim of sources 40–46 that is wrong and no claim asserted as proved without a proof. The four claims of source 43 that the intake left unchecked (finite-abelian orthogonality, the certification threshold $1 - 2^{1-d}$, the fourth-power triangle bound, the quartic value $11/16$) were checked at the write and hold in every test (note in Section F.48). Corrected or qualified in dated notes, with the first wording kept as delivered: credit, namely source 41’s reuse of source 32’s text (its Sections 2 and 5 verbatim) and of sources 26’s and 32’s bounds and method without citing either; source 45’s re-derivation of sources 28’s and 34’s single-pair results and of sources 29’s and 34’s moment bound without citation; source 46’s silence on source 36 and on source 43, and source 43’s on source 46; source 44’s citation of sources 25 and 27 by library identifiers; source 45’s interval “ $0.55743364 \leq c_4^{\text{odd}}$ ”, correct but not the best known lower end; the placement’s phrase “26’s cube-root length cap becomes a square root”, made precise (only the modulus cap $(N/8)^{1/3}$ becomes $\sqrt{N/2}$; note after source 42’s Theorem I.74.1); source 45’s description of this report at its pin as holding sources 01–14 (true then). Moved to questions as a marked question (not proved here): that source 40’s Corollary II.170.1 gives the formal `lemma_15_5` with an explicit threshold, merged with the same marked question for source 35 (item II.50). Kept as evidence only: the intake’s numerical minima of the U^4/U^2 ratio near 10.2571; source 44’s value $b_{2,8}$, which rests on its finite certificate (the method was checked independently only for $d \leq 4$); the numerical parts of the write’s

checks of source 43. Inputs used and not proved in the sources: the inverse Pollard theorem and Hegyvári's recurrence (source 45, credited); Tao–Ziegler's and Tidor's integration criterion (sources 43 and 46, credited); the Kolountzakis–Révész packing bound and Bhatia–Davis (source 42, proved in the needed form); Maynard's theorem for context only (source 41).

Answered or advanced inside the report. Source 26's question on the $(N/8)^{1/3}$ ceiling is answered by source 42 and its higher-degree question, with source 32's degree-three question, by source 41 (item I.19); source 27's two questions on $b_{2,d,p}$ and on intermediate characteristics are answered by source 44 (item II.38); source 14's Question 4 is advanced by source 45 to the range $\varepsilon < 1/9$ with all equality cases (item II.31); item III.1 is advanced to $0.5587725 \leq c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$ by source 45; source 36's characteristic-two question is answered for symmetric trilinear forms by sources 46 and 43, and its nonclassical and composite questions advanced by source 43 (item IV.9); the retention questions of sources 04 and 17 gain source 40's kernel-class barrier (item II.7).

Still pending. Sources 47–59 (placed by f8bc5e2ec and 7dc1ee29d while this write was in progress) are in this directory and not used in this text. Their placements record findings for items III.1 (source 49: $c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$), IV.10 (source 47: the quartic value $11/16$ for every n), II.53 (source 57: the third energy level), II.41 (source 50: a written counterexample to the catalogue lemma_16_10), II.46 (source 55: endpoint stability in every U^d) and others; the items below still state the questions as open.

[write] 7 October 2026; sources 47–53. *Sources 47–53.* They ask 57 further questions: source 47 nine (its Research questions in Section V.176, “47 Q1”–“47 Q9”), source 48 ten (Section V.181), source 49 sixteen (its Questions 10.1–10.16, Section V.186) and one question paragraph at the end of its Section 8 (Section II.235.11), source 50 ten (Section V.190), source 51 four (the enumerated items of Section V.195), source 52 four (paragraphs of Section V.198) and source 53 four (the enumerated items of Section V.202). They are added to the themes below as four new items or merged into existing items by dated sentences. The counts per theme become 72, 228, 120, 54 and 76 (550 questions; the headings keep the counts of sources 01–05).

Unproved or wrong claims in sources 47–53. The intake and the ninth write found no claim of sources 47–53 that is wrong and no claim asserted as proved without a proof; source 49's Section 8 is stated, and printed, as an implication from I1–I7, and sources 51 and 53 import their analytic inputs explicitly from proved Lean theorems at their pins. Two of the sources prove that repository propositions are false (source 50: lemma_16_10; source 52: two preserved Stage 13.7 encodings); those are findings about the encodings, recorded with their counterexamples in the sources' own theorems and in the crosswalk. Corrected or qualified in dated notes, with the first wording kept as delivered: credit, namely source 49's silence on source 45 (which postdates its pin; its interval's lower end is not the best known), the silence of sources 51, 52 and 53 on one another and on sources 26, 32, 41 and 42, source 53's repetition of source 51's lemmas, and the citation of sources 43 and 17 by library file names in sources 47 and 48; statements made false by later repository work and annotated at HEAD, with the pinned wording kept: source 53's “the Lean threshold is existential” and the absence of an explicit uniform threshold noted by source 51 (the repository has an explicit threshold since 93d436414), source 50's description of section16_local_affine_lift, source 52's location of lemma_13_7_holds, the transcription line numbers in source 49's crosswalk table, and the wrapper line ranges cited by source 53; and the eighth write's own description of source 48, from its placement, as a constructive form of the Fejér selections of sources 17 and 40 (source 48 treats source 17's restriction only; note in Section F.53). Moved to questions: none needed. Kept as evidence only: the intake's random-restart and exhaustive-form checks of source 47's thresholds, the intake's numerical U^4/U^2 minima near 10.2571, source 50's small word on $\mathbb{F}_{257}$, source 48's worked examples. Inputs used and not proved in the sources: I1–I7 (source 49, stated as hypotheses); the analytic Section 13 constructions (sources 51 and 53, imported from proved Lean theorems); nonclassical integration (source 47, credited to Tidor); Hoeffding's inequality (source 50, reproved for Bernoulli variables).

Answered or advanced inside the report. Source 43's quartic question is answered by source 47 (item IV.10); source 17's Question 9 in its ambient-size form by source 48 (item II.23); source 22's exactness question for $k \leq 3$ by source 49 (item II.42); item III.1 is advanced to $0.5587725 \leq c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ by source 49; the search of sources 19, 23, 28 and 39 for controls closing Lemma 16.10 is answered negatively for the catalogue form by source 50 (item II.41).

Still pending. Sources 54–59 (placed by 7dc1ee29d) are in this directory and not used in this text. Their placement records findings for items II.53 (source 57: the third energy level), II.46 (source 55: endpoint stability in every U^d), II.23 and the Section 9 material (source 58: Lemma 9.3 on every cyclic group) and others. Sources 60–69 (placed by 39d3295c4 and 06b30ca52 during this write) are not used either; their placement records findings for items III.1 (source 63: $\lim_d c_d^{\text{odd}} = 1/3$), II.49 (source 60: source 33’s Questions 2–3), II.8 (sources 61, 62 and 68: restriction on finite abelian groups and on every cyclic modulus), IV.11 (sources 65 and 69: source 47’s Questions 1–2), II.53 (sources 66 and 68) and others. The items below still state those questions as open.

[write] 7 October 2026; sources 54–59. *Sources 54–59.* They ask 63 further questions: source 54 eleven (its Questions 11.1–11.11, Section V.206), source 55 ten (Questions 12.1–12.10, Section V.208), source 56 ten (Questions 12.1–12.10, Section V.214), source 57 twelve (Research questions 1–12, Section V.216, “57 Q1”–“57 Q12”), source 58 ten (Questions 13.1–13.10, Section V.219) and source 59 ten (Questions 12.1–12.10, Section V.223). They are added to the themes below as four new items or merged into existing items by dated sentences. The counts per theme become 72, 281, 120, 54 and 86 (613 questions; the headings keep the counts of sources 01–05).

Unproved or wrong claims in sources 54–59. The intake and the tenth write found no claim of sources 54–59 that is wrong and no claim asserted as proved without a proof. Imported results are stated as such: source 55’s global entry is the published Eisner–Tao theorem; source 57 imports Gowers’s Proposition 6.1, Corollary 7.6 and Lemma 12.4 with their constants and the repository’s uniform square extraction at its quantifier order; source 59’s global affine consequences rest on the published correction theorem of Ben-Or, Coppersmith, Luby and Rubinfeld, its doubling corollary on Reiher–Schoen and Plünnecke–Ruzsa, its lonely-runner transfer on Bedert. One source reports an error elsewhere: source 59 shows that a displayed probability estimate in the published proof of Green and Tao’s Lemma 8.3 fails for an element of order five ($3/5$, not at most $1/2$); it is kept as a counterexample to that estimate only, as source 59 states it, not to the lemma or to the inverse theorem (note in Section II.252). Corrected or qualified in dated notes, with the first wording kept as delivered: credit, namely source 55’s silence on sources 31 and 38 and source 54’s on source 31 (both placed but unwritten at their pins), and the mutual silence of sources 48 and 58; source 57’s delivered `cleveref` names (“Theorem” for every statement on its counter); statements made stale by later repository work (source 57’s “written-proof gap”, its inexplicit N_{sq} , the integration suggestions of sources 54–57), each annotated at HEAD.

Answered or advanced inside the report. Source 45’s first two questions are answered by source 57 (item II.53); source 17’s composite-modulus question by source 58 for one balanced relation, and source 24’s mixed-torsion question by source 59 (item II.8); source 11’s question how far the radius $1/7$ extends is advanced by source 59 to $4/9$ and $1/2$ (item II.19), and its cutoff question to 16 for targets without involutions (item II.21); source 57’s threshold question (its Research question 8) is answered by source 53 (item II.55); the coefficient questions of sources 31 and 38 at the endpoint gain source 55’s all-degree coefficients (item II.46).

Still pending. Sources 60–64 and 65–69 were placed by 39d3295c4 and 06b30ca52 (“Place batch 126” and “Place batch 127”) before this write began, and sources 70–75, 76–81 and 82–88 by b788eda88, 487859db5 and bbde2b873 (“Place batch 128” to “Place batch 130”) while it was in progress; none of them is used in this text. Among the findings their placements record, those that bear on sources 54–59 are: source 61 finds source 59’s restriction exponent $(2m + 1) \log_2(5m)$ independently and gives an exact, threshold-free restriction on arbitrary finite abelian groups; sources 62 and 68 give Lemma 9.3 on every cyclic modulus by other methods, 68 with exponent 53 (source 58: 54); source 66 gives the endpoint partitions and the fourth energy level beside source 57’s Section 6, uncited (source 57 has the better third-level cutoff, $m \geq 7$); source 67 gives the non-affine energy maximum from $N \geq 13$, with 13 optimal for elementary abelian 3-group targets (source 59’s $\mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/3\mathbb{Z}$ example, uncited there); source 70’s headline is the printed Theorem 13.12, already proved in writing by source 57 and formally, with a near-maximal branch on every cyclic modulus; source 72 gives a near-maximal frequency map on every finite abelian group beside source 57’s Section 5; source 73 gives an explicit entry threshold $\varepsilon \leq 10^{-d}$ and radius $q^{-5/2}$ for source 55’s sharp expansion, which would answer source 55’s Question V.208.4; and sources 76–83 give sharp hereditary energy-rigidity constants for full-domain maps, to be printed after source 56 in the affine chapter. The items below state those questions as they stand after this write.

[write] **7 October 2026; sources 60–69.** *Sources 60–69.* They ask 104 further questions: source 60 fourteen (its Research questions 1–14, Section V.226, “60 Q1”–“60 Q14”), source 61 nine (Questions 12.1–12.9, Section V.230), source 62 eight (Questions 12.1–12.8, Section V.235), source 63 thirteen (Research questions 7.1–7.13, Section V.237), source 64 eight (Questions 10.1–10.8, Section V.240), source 65 nine (Research questions 1–9, Section V.244, “65 Q1”–“65 Q9”), source 66 ten (Conjectures 13.1 and 13.3 and Questions 13.2 and 13.4–13.10, Section V.249), source 67 eleven (Questions 7.1–7.11, Section V.254), source 68 twelve (Questions 8.2–8.13, Section V.256) and source 69 ten (Research questions 1–10, Section V.261, “69 Q1”–“69 Q10”). They are added to the themes below as one new item (III.31) or merged into existing items by dated sentences. The counts per theme become 74, 332, 141, 76 and 94 (717 questions; the headings keep the counts of sources 01–05).

Unproved or wrong claims in sources 60–69. The eleventh write, re-deriving the intake’s findings, found no claim of sources 60–69 that is wrong and no claim asserted as proved without a proof. One misprint is recorded with its correction: source 61’s list of seven numerators for the eight orders $q = 2, \dots, 9$ in the proof of its root-grid lemma (the eighth value is 3; the lemma holds as stated; note there). Imported results are stated as such: source 60’s connected limit uses Christ and Iliopoulou’s Riesz–Sobolev inequality; source 62’s refined envelope the Rosser–Schoenfeld totient estimate; source 64’s endpoint bounds the Shi–Dong colouring theorem, a 2026 preprint, and Deng–Tidor–Zhao’s realization (with its own prime-modulus proof); source 67 the Yang–Xie moment criterion and Weil’s bound; source 68’s uniform lower bound the Riesz–Sobolev inequality on the circle; the integration criterion of sources 65 and 69 is Tidor’s (reproved). Corrected or qualified in dated notes, with the first wording kept as delivered: credit, namely source 63’s silence on source 16 and its credit to a predecessor that was never delivered; source 64’s silence on sources 10 and 14 (its dimension barrier is source 10’s threshold); source 67’s unattributed use of source 59’s order-twelve example; and the same-day independent overlaps of sources 65 and 69, 66 and 57, 61 and 59, 62 and 58; statements made stale by later work (sources 60 and 63 quote source 45’s 4795/832, since improved by source 49; source 60’s norm bracket, superseded by source 63; the integration suggestions of sources 60, 61 and 64), each annotated.

Answered or advanced inside the report. The limit questions of sources 28 and 29 and source 60’s Research questions 1, 2 and the first part of 3 are answered by source 63 (item III.1); source 33’s Questions 2, 3 (equality) and 8 by source 60, and its Question 1 advanced (item II.49); the intake’s marked threshold composition is proved with one power of L less by source 60 (item II.6); source 45’s Questions 6.6–6.8 and source 14’s Question 4 (range and equality parts) by source 68 (item II.31); source 58’s Question 13.1 is given its certificate and its Question 13.2 advanced by source 68 (item II.7); source 57’s Research questions 2 (at the fourth level) and 3 (for hole models) by source 66, and its Research question 1 and source 45’s Questions 6.2–6.3 advanced (item II.53); source 11’s cutoff question is settled for elementary abelian 3-group targets by source 67 (item II.21); source 47’s Research questions 1, 2 and the tensor part of 4 by sources 65 and 69 (item IV.11); source 16’s Question 1 to first order by source 63 (item IV.7).

Still pending. Sources 70–75, 76–81 and 82–88 were placed by b788eda88, 487859db5 and bbde2b873 (“Place batch 128” to “Place batch 130”) before this write began, and sources 89–94, 95–96, 97–98 and 99–100 by 39beb65f2, 09f5b745a, 9599d8566 and 53eedc74c (“Place batch 131” to “Place batch 134”) while it was in progress; none of them is used in this text. Their placements record findings for these items: II.55 (source 70: the printed Theorem 13.12, already proved in writing by source 57 and formally, with a near-maximal branch on every cyclic modulus and a lower-range exponent conditional on Lau’s recurrence theorem; source 71: the Stage 13.7 obstruction of source 52 on fixed mixed prime support); II.2 and II.58 (source 72: a near-maximal frequency map on every finite abelian group beside source 57’s Section 5; source 73: an explicit entry $\varepsilon \leq 10^{-d}$ and radius $q^{-5/2}$ for source 55’s expansion, which would answer source 55’s Question V.208.4); II.41 and II.56 (sources 74 and 87: the packaged Lemma 16.10 premises and the contextual lift contract fail on sparse and positive-density slabs; source 75: a deletion fraction at parameter one); II.19 (sources 75–83 and 85: sharp hereditary relative-energy constants for full-domain maps, $3/4$, $19/27$, $\rho = (\sqrt{6} - 1)/(\sqrt{6} - 1 + \sqrt{\sqrt{6} - 2})$, and on $\mathbb{F}_3$ -spaces $5/8$, $53/81$, λ and $5/9$, to be printed after source 56 in the affine chapter); V.1 (sources 84 and 86: comparisons of the repository’s explicit five-term thresholds); II.32 and II.50 (source 88: product-property refinements of source 15’s chapter, with a selection corollary for source 35’s); I.20 (sources 90 and 93: the complete four-frequency missing-arc profile Φ_4 , independently, with grid corrections); II.23 and II.21 (source 93: affine recovery below graph defect $1/4$ for every abelian target, and the non-affine maximum $N^3 - 4N^2 + (10 + 2\kappa)N - (6 + 2\kappa)$ from $N \geq 13$, optimal, for every

target; source 90: a sharp bound for maps $C_N \rightarrow W$, $N \geq 12$); III.1 (source 93: $c_d(F) - D(F) = O(d2^{-d/2})$ uniformly over odd-order families, a rate that would supersede source 63's); IV.9 (sources 89, 91 and 92: the first nonclassical obstruction in odd characteristic, its repair subspaces and covers, and depth reduction); III.26 (source 94: sharp three-term counting with arbitrary homomorphisms, beside source 20's chapter); IV.10 (source 95: simultaneous integration of two Boolean cubics through the generic half-rank of their pencil); II.19, II.49 and III.1 (source 96: the affine cubic $4t - 10t^2 + 6t^3$ on the whole range $t \leq 2/3$ for every odd-order target; source 60's odd-target arrangement repair and the bracket $[15/256, 23/256]$, proved independently, in a version with no cubic term and no smallness condition; source 33's erasure-robust repair without its additive floor; $\|f\|_{U^4}^{16} \geq (843/100) \|f\|_{U^2}^{16}$, that is $c_4^{\text{odd}} \leq (100/843)^{1/4} = 0.5868716$); IV.11 (source 97: the maximizing functions of the canonical tensor for $d \geq 4$, $e(P + \theta|u|)$ with $\deg P \leq d - 1$, and dimension-free stability with optimal leading coefficient, which would answer the function questions of sources 65 and 69); II.19 (source 98: a is affine iff $r_3(a) > 5/8$, sharp, for full-domain maps between arbitrary abelian groups); II.48 (source 99: fixed-radius Bohr sets in every large prime cyclic group with exactly $2^r - 1$ runs, so $(2^r - 1)/2 \leq c_r^{\text{prime}}$, which would answer source 30's prime-order question and raise the lower rank base from $\sqrt{3}$ to 2); III.25 and III.8 (source 100: sets of density δ in every large prime cyclic group whose k -term count deviates by a nonzero constant times the squared U^d norm, so that no $o(u^2)$ counting error is possible, which would answer source 20's Questions 1 and 3). Among them, those bearing on sources 60–69: source 97 on the function questions of sources 65 and 69; source 96 on source 60's arrangement chapter (the same results, independently; source 60 arrived first and is printed in its form) and on source 67's characteristic-three range; source 93 on source 67's recovery threshold and cutoff (for every abelian target where source 67 has elementary abelian 3-groups) and on source 63's rate; sources 90 and 93 on source 67's Fourier profiles; sources 89, 91 and 92 on the odd-characteristic questions of sources 65 and 69; sources 80–83 work on $\mathbb{F}_3$ -spaces beside source 67's characteristic-three chapter (domains, where source 67 has targets); and source 74's and 87's slab obstructions concern the Lemma 16.10 interface whose affine count source 67 sharpens. The items below state those questions as they stand after this write.

Unproved or wrong claims. The intake and the write found no claim of sources 01–05 that is wrong, and no claim asserted as proved without a proof. Three inputs are used but not proved in the sources. The classical threshold-function counts (Zuev; Kahn–Komlós–Szemerédi; Baldi–Vershynin) enter source 01's asymptotics and are published theorems. The partition input $\text{PP}(K, R_0)$ of source 02 is now a formal theorem for polynomial phases on intervals (Section V.1). The “independent mathematical reviews” of source 03 are not shipped, and nothing here rests on them. Two values found at intake are not certified, $C_5 = 8410$ and $A_4 = 770$; they are item IV.2.

Answered inside the report (dated notes at the questions). Source 04's Question 8, the leading constant of $\log_2 F_k/k^2$ and $\log_2 A_k/k^2$, is answered by source 01: it is 1, using the classical threshold count, and source 01 certifies $F_4 = C_4 = 154$. Source 03's Research question on c_d^{odd} is answered for $d = 3$ by source 04's Theorem III.3.4: $c_3^{\text{odd}} = 2^{-1/2}$, attained by real parts of characters. With it, source 03's Remark III.2.9 is settled at $d = 3$. Source 04's Question 3 is the same question as source 03's for $d \geq 4$, and is merged with it (item III.1).

Theme I: density transfer and phase-flat partitions (14 questions)

- I.1 Higher-degree phase partitions** (01 Question 8; 05 paragraph 6). The density transfer is proved for any short-arc phase (source 01's Theorem I.2.1; source 05's Proposition I.17.2 and Corollary I.19.4). Missing: polynomial or nilsequence partitions with a quantitative minimum cell length and small occupied exceptional mass (01), and density-sensitive oscillation budgets tracked through geometric complexity with all losses parameterized (05). Simultaneous linear approximation does not give them, and Theorem I.18.5 does not transfer to nonlinear phases. [6 October 2026; sources 11–14] Source 13 (its Question 3) asks the same for degree $k \geq 2$: which hypotheses permit simultaneous partition exponents polynomial in the number of phases;

the coupled lower-degree coefficients after restriction are the obstacle, and quadratics sharing a leading coefficient are a suggested first case. [6 October 2026; sources 15–20] Source 13’s Question 3 is answered, conditionally on Maynard’s published recurrence theorem: source 18’s Theorem I.32.2 with Corollaries I.32.4 and I.32.5 (all quadratics: every $a < 1/(126q^2 + 129q + 2)$ with Lau’s numbers; a shared leading coefficient: $a < 1/(99q + 98)$ from the scalar Lemma 5.5 alone) and, independently, source 19’s Theorem I.34.3 (degree profiles, $O_k(q^k)$). What remains is the gap to the obstructions (item I.15).

- I.2** *The best transfer under arithmetic constraints* (01 Question 5; 05 paragraph 4). The extremizers of Theorem I.2.1 rotate cell arcs freely, and the disk radius and the capacity inequality are separately sharp (05). Asked: how much stronger the tail becomes for one fixed cyclic character on a common-step partition with nearly equal cells (01), and the exact feasible region of block positive mass, correlation, density and phase width (05). [6 October 2026; sources 11–14] Source 13 (its Question 7) asks which extremizers of its joint frontier (Theorem I.28.3) are ruled out when the phase is a polynomial on progression cells.
- I.3** *Finite partition profiles and the size-increment trade-off* (05 paragraph 1; 01 Question 6). The universal constant $1/3$ is settled (Theorem I.16.4). Asked: the exact best minimum length for each (r, s, M) and slope, and the feasible pairs of minimum and maximum lengths (05); the best uniform pair of length and increment, combining (I.3.4) with an optimal finite partition, including the final incomplete blocks (01). [6 October 2026; sources 06–10] Source 10 (its Question 3) asks the same for its phase partition: the best universal minimum block length for a prescribed phase-disc radius with all integer rounding retained, first at fixed η/δ as $\delta \rightarrow 0$; its density cap is sharp, its passage from one approximate period to blocks of lengths $L, \dots, 2L - 1$ may not be.
- I.4** *Retaining more than one density statistic* (01 Question 7). Theorem I.4.2 solves the optimization once stop-loss statistics are chosen. Missing: structural lemmas that propagate those statistics through repeated localization at acceptable cost.
- I.5** *Heterogeneous phase budgets* (02 Subsection 11.5; 03 Research question 5). The analytic constraint is exact (source 02’s heterogeneous inequality in Section I.8; source 03’s Lemma I.14.3). Missing: a refinement theorem whose cost depends on the oscillation allowed in each cell, so that budgets can be allocated by local mass. Any better $\text{PP}(K, R_0)$ transfers automatically to Theorem I.11.1. [6 October 2026; sources 11–14] Source 13 (its Question 6) asks for a joint frontier retaining all phase radii s_j , the cell weights and the measured $V = \mathbb{E}|f|$, by a dual optimization over supporting lines.
- I.6** *The finite-cell extremal problem* (02 Subsection 11.6). Theorem I.9.1 solves the weighted total-mass problem. Open: the largest good cell when M and the allowed cell sizes are fixed, with integer cardinalities, which would sharpen the $M - 1$ corollary. [6 October 2026; sources 06–10] Advanced by source 06: with M fixed and the cell masses free, its Theorem I.21.1 gives the exact optimum $\min\{P_+/(b(M - 1)), S/(M - 2)\}$ with explicit extremizers (note after Corollary I.9.3). Still open: prescribed allowed cell sizes and integer cardinalities at a fixed N . [6 October 2026; sources 11–14] Source 13 asks the same (its Question 5: the largest guaranteed high-average cell for a fixed number of cells); in the weighted model with M fixed and the cell masses free it is answered by source 06’s Theorem I.21.1.
- I.7** *Constants in the simultaneous minimax theorem* (05 paragraph 2). The exponent $1/(m + 1)$ and the dependence on ε are sharp; the upper constant c_m is poor. Asked: the optimal leading

constants, first for one or two phases (source 01's Theorem I.7.2 improves the one-phase constant), the effect of a common difference, and the order of limits. [6 October 2026; sources 11–14] Source 13 (its Question 2) asks for the best leading constants of its flexible partition theorem for fixed q and anisotropic tolerances; its constructive and probabilistic constants differ by less than a factor 4.

- I.8** *Anisotropic lower bounds* (05 paragraph 3). Is there a matching minimax upper bound in all anisotropic regimes, with the regime decomposition that loose tolerances require? [6 October 2026; sources 11–14] *Advanced by source 13*: its Theorem I.26.5 (random slopes) with Corollary I.26.3 gives the order $(N \prod_i \varepsilon_i)^{1/(m+1)}$ for independent tolerances $0 < \varepsilon_i < 1/2$, constants within a factor below 4. Loose tolerances $\varepsilon_i \geq 1/2$, where the regime decomposition is needed, are not covered.
- I.9** *Adaptive selection of frequencies and tolerances* (05 paragraph 5). Proposition I.19.3 solves the continuous allocation. Missing: subset selection and the exact ceiling problem, relative to certified Fourier inputs.
- I.10** *Interval-localized spectral extraction* (05 paragraph 8). Corollary I.19.4 corrects for the Gram matrix by its top eigenvalue only. Missing: certified conditioning bounds for specific frequency configurations, using more of the Gram matrix. [6 October 2026; sources 06–10] Source 10 (its Question 4) asks the companion problem: when several approximately constant characters or polynomial phases live on one partition, the best Gram-matrix lower bound for the conditional second moment, to be combined with its variance cap (I.23.3); the common partition has to be supplied by the geometry, not assumed.

Added in the second write, 6 October 2026 (sources 06–10):

- I.11** *Using the mass of all dense cells* (06 Question 8). Source 06's Theorem I.21.1 controls the total mass above every threshold, not only one cell, and its extremizers show that the scalar data cannot give more. Asked: whether a branching density-increment argument, keeping geometric information about the collection of high cells, improves a structural or counting bound. Related to item I.4.

Added in the third write, 6 October 2026 (sources 11–14):

- I.12** *Prescribed lengths at the flexible frontier* (13 Question 1). Can a complete partition into lengths $v - 1$ or v reach $(q + 1)a + qb = 1$, or is there an arithmetic obstruction between it and the proved region $(q + 2)a + qb \leq 1$ (Corollary I.26.9)? The loss comes from the conductor $(v - 1)(v - 2)$; with an exceptional set of mass $o(1)$ the flexible exponent is reached. [6 October 2026; sources 15–20] Source 19 asks the same for nonlinear phases (its Question 8 and S6(4): prescribed lengths without the square-length loss of its Lemma I.34.4); *partly answered by source 18*, whose Theorem I.32.2 pays the square once, at the base (note after Corollary I.32.4). Source 18 quantifies the cost of prescribed lengths in degree one (its exponent recursion), without answering source 13's question. [6 October 2026; sources 21–27] *Answered by source 23*: there is an arithmetic obstruction. No partition into lengths $v - 1, v$ reaches beyond $(q + 2)a + qb \leq 1$, even with cell-dependent differences (Proposition I.38.1; the lower example is realizable on a prime cyclic group), so the exponent $1/(3q + 2)$ at Gowers's $b = 2a$ is optimal; for lengths $L, L + 1$ the threshold is $L^{q+2} / \prod_i \varepsilon_i (1 + o(1))$ (Corollary I.37.1). The nonlinear prescribed-length question of source 19 stays open (item I.15).

- I.13** *Bounded-height presentations* (13 Question 4). Choose a generating family efficiently so that (I.26.6) is close to optimal for a given collection of modular phases, with certified generator matrices and budgets.

Added in the fourth write, 6 October 2026 (sources 15–20):

- I.14** *Partitions controlling their own steps* (16 Questions 5–7). Source 16’s sufficient exponent for equal powers is of order q^{-2} (Theorem I.30.5, conditional on Lau’s preprint), its obstruction of order q^{-1} (Corollary I.31.2): can a universal theorem reach c/q ; is the strict interior of $(q+1)a + 2\sum_i b_i \leq 1$ sufficient; can a weighted extraction of one progression beat the full-partition obstruction at the Lemma 13.5 interface?
- I.15** *The nonlinear partition exponent gap* (18 Question 9; 19 Question 7 and S6(1)–(2)). For fixed $k \geq 2$, the sufficient frontiers (source 18’s $126q^2 + 129q + 2$ for quadratics, source 19’s $O_k(q^k)$) against the necessary $(qk(k+1)/2 + 1)a + kqb \leq 1$, equivalently $(S+1)a + Db \leq 1$ (for quadratics with $b = 2a$: $a \leq 1/(7q+1)$): reduce the sufficient denominator to order q , or attain the necessary region; decide first whether flexible lengths and lengths $v-1, v$ have the same frontier (source 18). [6 October 2026; sources 21–27] Source 23 (its Question 6) asks for the prescribed-length obstruction for polynomial phases: a higher-degree analogue of its mixed-radix construction, with a matching necessity theorem. [6 October 2026; sources 34–39] Advanced by source 35 for prescribed lengths: explicit phases with one prime denominator per degree force $D \mid h$ for every flat progression, the exact thresholds are $DC(S)$, and the necessary frontier for lengths L or $L+1$ is $(W+2)a + \sum_i d_i b_i \leq 1$ (quadratics at $b = 2a$: $1/(7\nu+2)$; Theorem I.61.2, Corollary I.61.6); this answers the construction half of source 23’s Question 6. The flexible-length obstruction of source 19 is unchanged. Source 35 asks (its Questions 10 and 12): a matching quadratic upper theorem, threshold $C_\nu L^{3\nu+2} \varepsilon^{-2\nu}$ up to subpower factors; and the same powers for prime moduli comparable with the interval length.
- I.16** *Effective dimension of the phases* (16 Question 8; 18 Question 10; 19 S6(5)). Replace the number of phases by an effective dimension: moving slopes in a low-dimensional bounded-height span (16), bounded-height rational dependencies among leading coefficients (19), and a recurrence formulation that also controls the coefficients created by translation, with the shared-leading family as a test case (18).
- I.17** *Numerical recurrence thresholds and coefficient-rank profiles* (16 Question 9; 18 Question 11; 19 Question 9 and S6(3)). Make the thresholds $X_0(C, q)$ and the constants $U_{j,m}$ of the cited recurrence theorems numerical, or formalize certificate-producing finite algorithms; give each degree its own bounded integer span and compute the resulting recurrence costs. An exponent with an uncontrolled starting threshold supports no finite comparison.

Added in the fifth write, 6 October 2026 (sources 21–27):

- I.18** *Prescribed-length thresholds beyond the leading term* (23 Questions 1 and 3–5). The second asymptotic term of $T_{\{L, L+1\}}(\varepsilon) = L^{q+2} / \prod_i \varepsilon_i (1 + O_q(1/L))$ (start with $q = 1$ and a rational tolerance); stability of the obstruction (must nearly worst phase vectors resemble nested rational quotients?); the leading constant when the allowed lengths span a fixed fraction θL of the minimum; the optimal exceptional-set profile at fixed η and for a general numerical semigroup (Theorems I.36.1, I.39.1). [6 October 2026; sources 34–39] Source 35 (its Question 11) asks how the threshold changes for the allowed lengths $L, \dots, L+w$ (its conductor $L[(L-1)/w]$) and how much an exceptional set reduces it.

I.19 *Corollary 5.8 without a scale threshold* (26 Questions 1–4). Relax the $(N/8)^{1/3}$ ceiling of Theorem I.41.1 (it comes from the sufficient condition $N \geq L^3$); an α -independent square-root guarantee; the optimal constants of the local-cover interface; and a threshold-free replacement for the phase-flat refinement in higher degree, with its small-correlation regime. The placement of source 32 records it as treating degree two; that of source 41 records a written proof of the whole statement in every degree, gain $\alpha/3$ at length $\max\{1, (N/M)^{1/D_k}/8\}$, and that of source 42 the square-root cap $\sqrt{N/2}$ in place of $(N/8)^{1/3}$. These are for the writes that add them; until a write checks them, the item stays open. [6 October 2026; sources 28–33] *Answered at degree two by source 32*: gain $\alpha/3$ at length $\max\{1, (N/M)^{1/39}/17\}$ over every cyclic group with no scale threshold (Theorem F.37.1); with source 26 the instances $k \leq 2$ of the formal statement have written proofs. Source 32’s own questions: a better explicit recurrence ($p \leq CR^\kappa$ would give the length exponent $1/(7\kappa + 4)$; Remark I.45.4 does not supply $\kappa = 4$), a more economical quadratic partition than power 38, parent-contained small-correlation increments, and degree three (32 Questions 1–4). Degree three and beyond stay open here; source 41’s placement records a proof in every degree, for its write. [7 October 2026; sources 40–46] *Answered in every degree by source 41*: gain $\alpha/3$ at length $\max\{1, (N/M)^{1/D_k}/8\}$, $D_1 = 4$, $D_k = K_k/4 + 1$, over every cyclic group with no scale premise (Theorem F.46.1), with the cubic bound $(N/M)^{1/4096}/8$ (Corollary I.68.5); so every instance of the formal statement has a written proof; the formal statement is now exact as well, by the repository’s own proof along the same architecture (`corollary_5_8_holds`; Section V.1). *The $(N/8)^{1/3}$ ceiling is answered by source 42*: it becomes $\sqrt{N/2}$ (Theorem I.74.1). Source 41 asks (its items 1–4): a fully explicit improved recurrence at all small scales, inserted into the exact recursion; a cheaper localization (e_k, c_k) ; a sharper final selection, trading the gain $\alpha/3$ against length with singleton and boundary cases kept; parent containment for the low-correlation branch and rectification to ordinary integer progressions. Source 32’s degree-three question is answered (source 41’s cubic bound); its other questions stand.

Added in the seventh write, 6 October 2026 (sources 34–39):

I.20 *Missing intervals beyond the first window* (37 Questions 1–6, 8 and 10). The minimax $\inf \max_{r \leq R} |\hat{\mu}(r)|$ when $a > 2\pi/(R + 1)$, with extremizers and transition points (a piecewise algebraic description in $\cos a$?); exact finite-grid corrections on $\mathbb{Z}/N\mathbb{Z}$, open against half-open exclusions; uniform stability constants as $R \rightarrow \infty$; unequal Fourier budgets ε_r and the region where the certificate of Section I.51 stays optimal; several forbidden arcs; the best constant of the recurrence conversion with the exact frontier $\rho_R(2\pi\delta)$; whether the structure of polynomial empirical measures $T^{-1} \sum_{n \leq T} \delta_{\alpha n^k}$ improves the recurrence; and which parts of the formula for ρ_R and its extremizing quadrature are classical in equivalent language. Source 37 is the first source of this report on Lemma 5.2 (Section I.49). [7 October 2026; sources 60–69] *Advanced by source 67*: the exact value near the equal-budget branch under unequal budgets, the certificate polytope and the Hessian rank (Theorems I.76.2–I.76.6). Source 67 asks (its Questions 7.1–7.2): $V_{R,a}(\varepsilon)$ throughout the nonnegative budget orthant, with the regions of constant extremal type; the subsequential limits of $N^2(\rho_{R,N}(a) - \rho_R(a_N))$ on grids.

Added in the eighth write, 7 October 2026 (sources 40–46):

I.21 *The arithmetic variance increment* (42 items 1–4). The best uniform coefficient under $N \geq CL^2$ for densities bounded away from 0 and 1, including $C = 1$ (Theorem I.73.1 gives a positive endpoint coefficient, not shown optimal); for fixed N , L and $|A|$, when the arithmetic bound is attained; the best supports and positive-definite weights in general finite abelian groups;

comparable gains inside a prescribed parent or cover, and nonwrapping integer conclusions (merged with source 41’s containment question, item I.19). Section I.62.

The ten items carry the fourteen questions: 01 Questions 5–8, 02 Subsections 11.5–11.6, 03 Research question 5, and 05 paragraphs 1–6 and 8. [6 October 2026; sources 06–10] With sources 06–10 the eleven items carry seventeen questions: also 06 Question 8 and 10 Questions 3–4. [6 October 2026; sources 11–14] With sources 11–14 the thirteen items carry twenty-four questions: also 13 Questions 1–7. [6 October 2026; sources 15–20] With sources 15–20 the seventeen items carry forty questions: also 16 Questions 5–9, 18 Questions 9–11, 19 Questions 7–9 and its five questions S6(1)–(5). [6 October 2026; sources 21–27] With sources 21–27 the nineteen items carry forty-nine questions: also 23 Questions 1 and 3–6 and 26 Questions 1–4. [6 October 2026; sources 28–33] With sources 28–33 the nineteen items carry fifty-three questions: also 32 Questions 1–4. [6 October 2026; sources 34–39] With sources 34–39 the twenty items carry sixty-four questions: also 35 Questions 10–12 and 37 Questions 1–6, 8 and 10. [7 October 2026; sources 40–46] With sources 40–46 the twenty-one items carry seventy-two questions: also 41 items 1–4 and 42 items 1–4. [7 October 2026; sources 60–69] With sources 60–69 the twenty-one items carry seventy-four questions: also 67 Questions 7.1–7.2.

Theme II: the inverse step (9 questions)

- II.1** *A correlation-preserving structural theorem for multivalued spectra* (02 Subsection 11.1). Proposition II.4.1 keeps all large derivative frequencies. Missing: an energy-to-structure step that keeps a fixed fraction of t_τ and exploits κ , controlling both losses. [6 October 2026; sources 28–33] *Answered in a prime-cyclic, bounded-fibre form by source 31*: its Theorem II.117.1 keeps an explicit fraction of the selected derivative correlation through an exact Freiman restriction on $\mathbb{F}_p$. Open: compatible restrictions across slices (31 Question 9), which is what Gowers’s later multilinear stages need.
- II.2** *Near equality for the coupled pair (f, w)* (02 Subsection 11.2). Proposition II.4.2 gives two alignments, and the equality examples show distinct geometries. Asked: a dimension-free description of f and w jointly when Theorem II.2.1 is nearly an equality; results on near-maximal Gowers norms (Eisner–Tao) are relevant but do not classify the weight. [6 October 2026; sources 28–33] *Advanced by sources 31 and 29*. Source 31 gives the exact optimal profile for the function (Theorem II.107.4), sharp selected-graph stability (Theorem II.108.1), a weight-defect inequality in the U^2 and coset topology and counterexamples to normalized L^1 control of the weights (Section II.109); source 29 the endpoint with soft and mixed selectors (Sections II.111–II.112). Still open: the unrestricted coupled near-equality, with the optimal coupled weight defect (31 Questions 10–11), optimal constants and a larger defect range of the joint theorem, the row constraint, and higher-degree joint selectors (29 Questions 5–7). [7 October 2026; sources 54–59] *Explicit local constants by sources 55 and 57*: source 55’s a posteriori local theorem (Theorem II.242.1) and source 57’s near-maximal extraction with constants 332 , $216s + 8$ and the cubic bound $\sqrt{1 - 6194\varepsilon} - \varepsilon$ (Theorems II.247.5, II.247.7, Corollary II.247.10). Also asked: an explicit Eisner–Tao entry threshold e_d , singly exponential in d (55 Question 12.4); sharper correction constants, the best ratio of squared cubic error to the Q_4 deficit (at least $1/8$), and better degree and characteristic dependence (57 Research questions 9–10).
- II.3** *Carrying profiles through the proof* (02 Subsection 11.3; 04 Q7). Keep the distribution of $m(\Delta_{\mathbf{h}}f)$ and $\mathcal{U}(\Delta_{\mathbf{h}}f)$ across slices instead of their supremum (02), and the cross-section

profile $\sigma = \mathbb{E}_r \beta_r^3$ through later restrictions (04). The difficulty in both: compatibility of structure between slices, and which operations preserve the profile. [6 October 2026; sources 11–14] Advanced by source 14: its Theorems II.44.1 and II.44.4 classify equality in the profile bound and its Corollary II.44.3 improves the denominator from cross-section sizes alone; propagation through later restrictions is still open. [6 October 2026; sources 15–20] Source 15 (its Question 6) asks when one can pass to a balanced subdomain while keeping enough arrangements to use ρ_3^{-6} ; its Example II.48.4 shows the cost of assuming balance.

- II.4** *The sharp fixed-threshold DRC profile* (03 Research question 3). The order of the best guaranteed $|K|/n$ under Lemma 7.4’s exact criterion as $\delta \rightarrow 0$, and whether the exponent 3 of Theorem II.5.2 can be improved using constraints of actual intersection matrices; the scalar argument only rules out a two-witness first-moment certificate. [6 October 2026; sources 21–27] Also asked by source 24 (its Questions 1–4: the dense-intersection extremal function $f(\delta)$, whether actual codegree geometry rescues two samples, profile-sensitive sampling with a verifiable certificate, and higher-arity scalar envelopes). Not advanced: source 24’s scalar facts are source 03’s (note after its Theorem II.75.1).
- II.5** *Optimizing the elementary extraction* (03 Research question 4). Optimize $\gamma, \varepsilon, \beta$, walk lengths and popularity thresholds in Proposition II.6.2, judging the pair of conclusions together and keeping the comparison with Reiher–Schoen explicit. [6 October 2026; sources 21–27] Source 24 (its Question 8) asks the same end to end, with the actual energy spectrum; its popularity-threshold optimizations (Sections II.77 and II.80) advance it for its own route.
- II.6** *Lowering the size threshold of the restriction theorem* (04 Q4), from $|B| \geq \Lambda_d(\alpha\eta)^{-2d}$, by reflecting the additive structure of B in the exceptional relations or by a better kernel, with repeated entries handled. [6 October 2026; sources 15–20] Also asked by source 17 (its Question 2: lower the degree $2m - 1$ in $(\alpha\eta)^{-1}$ of its threshold, keeping its retention powers; its Question 3: the extremal filter maximizing $\sum_n |\hat{h}(n)|^m$ at prescribed mean, the kernel half of this question). Source 04’s weighted count already has degree m and β^{-1} (note after Theorem II.61.2). *Marked question (the intake’s composition, in neither source; not proved at the write)*: source 17’s Fejér kernel with source 04’s per-tuple lower bound and its weighted rank-two count of exceptional tuples (Lemma II.9.3) should keep source 17’s retention $1/(2\eta L^m)$ under the threshold $|B| \geq 2(1 + \binom{m}{2})L^m$, about $2^{202.3}$ at the Corollary 9.4 point $\gamma = \eta = \beta = 1/2$, below both sources’ thresholds; the intake checked the two needed inequalities in twelve small instances. Missing: a written proof that the per-tuple lower bound and the weighted count combine without a D_G loss for repeated and degenerate tuples. [6 October 2026; sources 28–33] Answered for fixed fields by source 29: threshold $8qS_{q,s}/(\alpha\eta)$, linear in $(\alpha\eta)^{-1}$, with every degenerate tuple kept (Theorem II.120.4), and none at a smaller coefficient (Corollary II.120.5); this is also source 17’s Question 2 in its fixed-characteristic form. Open: the prime-cyclic and growing-modulus case (item II.8). Source 29 asks for a bound by the actual affine-span statistics of B in place of $z \prod_i (z + q^i)$ and an efficient deterministic restriction (29 Questions 10–11). [7 October 2026; sources 54–59] Source 58 gives the explicit all-moduli threshold $|B| \geq (4/(\alpha\eta))^{54}$ (Theorem A) and asks (its Question 13.4) for the best point-count threshold for a prescribed retention, using the collision and torsion profiles. [7 October 2026; sources 60–69] *The marked composition above is proved by source 60 with a smaller threshold*: on prime moduli $b \geq 2(1 + \binom{m}{2})L^{m-1}$ keeps source 17’s retention (Corollary II.330.5), one power of L less than marked. Source 60 asks (its Research questions 10 and 12): the best translated positive kernel for the signed objective $\eta EX - \mathbb{E}Y$; whether the exponent $m - 1$ is necessary uniformly in the prime.

- II.7** *The optimal retention law* (04 Q5). Proposition II.9.6 limits one expectation budget only. Asked: upper-bound examples or a better selection method, and simultaneous guarantees on $|B'|$, its energies and the respected density. [6 October 2026; sources 15–20] Source 17 (its Question 1) asks the same for prime cyclic groups: whether losses of order $\alpha^m \eta^{m-1}$ in the tuple count are needed; its Proposition II.62.2 settles only the scalar certificate. [7 October 2026; sources 40–46] *Source 40’s barrier*: within homogeneous positive-Fourier kernels, gain Q forces good survival at most $Q^{-(s-1)}$, attained by Fejér kernels (Theorem II.172.2, Proposition II.172.3). Source 40 asks (its Question 7) whether vertex-dependent kernels or adaptive selection beat the scale $\alpha^s \eta^{s-1}$ with a usable exceptional-set estimate. [7 October 2026; sources 54–59] *Source 58’s barrier*: within normalized positive Fourier filters the leading principal cost is exactly $m \log m / (2 \log a)$ (Theorem B), the multi-layer counterpart of source 40’s barrier. Source 58 asks (its Questions 13.1–13.3): the sixteen-variable optimum $I_{16,2}$, below 52; the next asymptotic term; a retention exponent $O(m)$ on every cyclic group by a method outside independent positive filters. [7 October 2026; sources 60–69] *Source 58’s Question 13.1 is given its certificate by source 68* (a degree-twelve filter of principal cost $51.445 < 52$, exponent 53) and its Question 13.2 advanced (two-parameter bounds, $\mathcal{I}_{m,a}/m \rightarrow 1 + 1/(2c)$); source 62 retains $\alpha^{31} \eta^{30}$, and $\alpha^{16} \eta^{15}$ up to a double logarithm that is necessary for its detector. Asked: a detector with no double logarithm (62 Question 12.1); $\sup_L \kappa_{m,L}^*$ (62 Question 12.2); whether $\alpha^m \eta^{m-1}$ is optimal for the theorem itself (62 Question 12.6); $\mathcal{I}_{16,2}$ and an exponent below 53 (68 Question 8.7); the second-order principal cost (68 Question 8.8); degree, principal and finite-size costs together (68 Question 8.9); filters for an observed order profile (68 Question 8.10); retention $O(m)$ beyond positive filters (68 Question 8.11, source 58’s Question 13.3).
- II.8** *Restriction beyond prime cyclic groups* (04 Q6): versions for $\mathbb{Z}/N\mathbb{Z}$, $\mathbb{F}_q^n$ and other codomains, with hypotheses that quantify kernels of the coefficient maps. [6 October 2026; sources 15–20] Source 17 (its Question 4) asks for composite cyclic moduli with explicit torsion dependence (Smith normal forms or gcds of coefficient minors); its fixed-characteristic Section II.63 answers the $\mathbb{F}_p$ -vector-space and weighted parts of this item, not $\mathbb{Z}/N\mathbb{Z}$. The formal all- N proof uses degree-one digit filters with fibres of size at most two (`Proofs15Restriction.lean`, lines 2083–2086, at `0b4890a1c`). *Marked question (the intake’s observation, in neither source; not proved at the write)*: source 17’s Lemmas II.60.3–II.60.5 should carry over verbatim when every prime factor of N exceeds $2L - 2$, since the relevant minors are integers of absolute value at most $2L - 2$ and every $0 < |t| < L$ is then a unit (the intake checked exact averages at $N = 77, 85$; with a factor at most $2L - 2$, unrespected tuples survived at about half the respected rate). Missing: the written proof, and the case of small prime factors. [6 October 2026; sources 21–27] Source 24 (its Question 7) asks for general finite abelian codomains with explicit dependence on the invariant factors; its Theorem II.79.5 settles the elementary abelian case of the Freiman extraction, independently of the dimension. [6 October 2026; sources 28–33] Source 29 (its Question 12) asks for a growing-modulus analogue of its fixed-field theorem, with constants useful as the prime grows, and eventually all cyclic moduli. [7 October 2026; sources 54–59] *Source 17’s Question 4 is answered for one balanced relation by source 58*: restriction on every cyclic modulus with explicit domain-torsion cost $\tau_G(2D)$ and repeated entries, exponent 54 for sixteen entries, every finite abelian target (Theorems A and II.282.1); the intake’s marked prime-factor case is contained in it. *Source 24’s Question 7 is answered by source 59*: target-independent restriction $C^{-(2k+1) \log_2(5k)}$ over arbitrary abelian targets (Corollary II.265.2). Source 58 asks (its Questions 13.5–13.7): low-exponent domains such as $(\mathbb{Z}/2\mathbb{Z})^s$; systems of relations with every collision pattern, first two independent relations;

filters adapted to the distribution of discrepancy orders. [7 October 2026; sources 60–69] *Further all-modulus theorems* (Section II.310): exponent 53 for arbitrary abelian targets (source 68), unit-dilation restriction for $\mathbb{Z}/N\mathbb{Z} \rightarrow \mathbb{Z}/M\mathbb{Z}$ (source 62), threshold-free exact restriction on arbitrary finite abelian groups (source 61). Asked: restriction with persistent short torsion through several translated selectors (60 Research question 11); a separating family for noncyclic targets without a bad rank dependence (62 Question 12.5); several relations with decaying collision errors through Smith invariants (68 Question 8.12, with source 58's Question 13.6).

Added in the second write, 6 October 2026 (sources 06–10):

- II.9** *The optimal nonuniform transport threshold* (07 Question 1). For fixed $k \geq 2$ and fibre ratio $\kappa > 1$, the largest universal δ forcing order- k relations, and the gain from a separate boundary bound. Proved: $1/(2[1 + (k - 1)\kappa])$ (Corollary II.15.3); sharp only at $\kappa = 1$ (Theorem II.16.1). Suggested: a finite search over fibre sizes and actual value assignments, not independent error profiles. [6 October 2026; sources 34–39] *Answered in source 07's model by source 34*: $1/(2 + (2k - 2)\kappa)$, which is source 07's sufficient threshold, is sharp for every $\kappa \geq 1$, already with zero boundary loss, so a separate boundary bound gives no improvement (Theorem II.163.1 and its corollary). The prime-order interface of the formal `lemma_10_9` is a different question.
- II.10** *The largest guaranteed Freiman order at Lemma 10.9* (07 Question 2), under the exact Lean hypotheses (order five proved) and under the ambient package (order nine proved). Five and nine are lower bounds; counterexamples must satisfy the nested local model, not only a scalar defect bound. [6 October 2026; sources 11–14] Source 12 (its Question 9) asks whether $3\varepsilon + r\kappa < 1$ is the correct universal scale for forcing order r on C from order two on S (its Theorem II.35.6), and which relation-complex parameters replace r .
- II.11** *Stability of winding obstructions* (07 Question 3): must a violated relation with total budget close to one come from a few cuts or carry boundaries, after a structural reduction?
- II.12** *Efficient certificates for many relations* (07 Question 4): compressing the root-law linear programs of Theorem II.18.3 over a large family of relations, for instance by translation symmetry or a presentation of the generated group; one relation is a finite linear program, exponentially many are not efficient.
- II.13** *Joint optimization of extraction and transport* (07 Question 5): building the regular component and its source law together, so that the domain is large and the relevant profiles small, accounting for lost cardinality and the relation between energy weights and the chosen law.
- II.14** *Higher-dimensional coherence* (07 Question 6): shared-vertex sampling on cubes or other diagrams in place of nested exceptional sets in bilinear and multilinear exactification; the joint law must be specified first and every marginal proved.
- II.15** *The earlier exponent $\eta^{1/5}$* (07 Question 7): which losses before the local model are necessary, and whether the fifth root can be improved while keeping a large regular component and direction set. Source 07 removes the later square root only.
- II.16** *A sharp vertical-fibre restriction theorem* (10 Question 5 and the question paragraph at the end of Section II.21). How close is $|B|/(k(h_k - 1) + 1)$ to the best bound determined by the defect fibre alone; can the factor k be reduced under additive structure of H_k ; can $|H_k|$ be bounded from a vertical energy profile, how does it grow with k , and when does an order- k

homomorphism extend to order $k + 1$ with small loss? Settled only for $H_k = \{0\}$. [6 October 2026; sources 21–27] Also asked by source 24 (its Questions 5–6: the smallest universal vertical-fibre exponent in place of $2k + 1$, and the largest set P with $(kP - kP) \cap F = \{0\}$ for a given symmetric defect set). [6 October 2026; sources 28–33] Also asked by source 31 (its Questions 5 and 8: optimize the retention using the set V_k itself, and replace the worst fibre size K^{2k+2} by a fibre profile). Not advanced: source 31’s restriction is source 10’s (note after its Theorem II.116.1).

II.17 *Overlaps in the bad-step union* (10 Question 6): Theorem II.21.3 uses a union bound over $E_j = \{d \neq 0 : jd \in F\}$; an exact treatment of their intersections, for example for cyclic prime-power codomains via valuations, would improve the general-codomain retention. [6 October 2026; sources 21–27] For codomains $\mathbb{F}_p^d$ the union over j is avoided altogether by source 24’s linear hashing (Theorem II.79.5), which beats the budget here whenever $Q_k \leq p^{d-1}$ and $k \geq 2$ (note there). [6 October 2026; sources 28–33] Source 31 (its Question 6) asks for its weighted restriction over composite cyclic and general finite abelian targets, charging the kernels of $1, \dots, kL$.

II.18 *Propagating a modern energy input* (10 Question 7): propagate the Reiher–Schoen retention (almost $\sqrt{\gamma}$, against $\gamma/4$ here) with the exact vertical fibre through the later Bohr-neighbourhood and local-linearity stages, producing a complete chain of inequalities rather than an isolated replacement of the exponent 1164. Related to item II.5. [6 October 2026; sources 21–27] Source 24 (its Question 9) asks the same for its Section 7 interfaces: which later lemmas survive the substitution, with every threshold rechecked. [6 October 2026; sources 28–33] Advanced by source 31: its Theorem II.118.1 is a weighted square-root mass theorem from Reiher–Schoen; it asks (Question 7) to use it, or a sharper weighted BSG theorem, to improve the powers of its weighted extraction.

Added in the third write, 6 October 2026 (sources 11–14):

II.19 *The affine-rigidity envelope and the next energy levels* (11 Questions 1–2). For fixed G, W and attainable t , $\max\{\mathcal{E}(f)/|G|^3 : d(f) = t\}$; how far the radii $1/7$ and $1/3$ extend (a joint optimization of two- and three-error patterns may help); and the next largest graph energies after one-point modifications, first for $\mathbb{F}_p \rightarrow \mathbb{F}_p$. [7 October 2026; sources 54–59] The radius $1/7$ is extended by source 59 to $4/9$, and to $1/2$ without 3-torsion, with coset equality and both endpoints resolved (Theorem II.250.4, Corollary II.250.6). Source 59 asks (its Questions 12.6–12.8 and 12.10): whether $\varepsilon \geq \Psi(t)$ holds on all of $t < 1/2$ (involution-free) or $t < 2/3$ (no 2- or 3-torsion) from a specified reference; the optimal support-colour frontier (α, β) for $t < 1/3$; a larger range for $\Psi_{\text{inv}}(t) = 4t - 12t^2 + 8t^3$ when the target has involutions; the sharp small- t envelope of higher graph energies E_r , $r \geq 3$. [7 October 2026; sources 60–69] Advanced by source 67: for elementary abelian 3-group targets $\varepsilon \geq \Psi(t)$ on the whole range $t \leq 2/3$, with remainders and all equality cases (Theorem II.294.1). Source 67 asks (its Questions 7.3–7.5): the largest target class with the full range up to $2/3$ and its equality cases; sharp stability near the coset extremizers; the attainable spectrum of (affine distance, defect) pairs.

II.20 *Near-extremizers at the right scale* (11 Question 3). Answered in substance at the write by Remark W3: for $\theta = \eta/(6t^3) \leq 1/100$ (odd targets) the error support is within $\tau(\theta)|S|$ of a unique coset and has one value off a fraction $\eta/(2t^2(1 - 7t))$ of it. Open: $\theta > 1/100$, and sharpness of the value term.

- II.21** *The exact cutoff for general groups* (11 Question 4). Remark W2: the cutoff 34 cannot simply be removed; the exceptions found have $N \leq 12$, with few homomorphisms or a non-affine quotient of small order, and the formula holds for $\mathbb{Z}/10, \mathbb{Z}/14, \mathbb{Z}/15, \mathbb{Z}/16 \rightarrow \mathbb{Z}/2$. *Marked question (the intake's derivation, in neither source; arithmetic checked at the write)*: replacing the step $\sum p_w^2 \leq \max p_w$ in Lemma II.28.1 by source 12's Lemma II.36.1 gives disagreement at most $\rho/(1 - \chi(2\rho))$; the radius $1/7$ is then reached at $\varepsilon = 6/49$ (solve $\sqrt{1 - 4\varepsilon} = 14\varepsilon - 1$; there $3\chi(2\varepsilon) = 3/7 < 1$), and since $\Psi_0(1/30) = 551/4500 < 6/49 < \Psi_0(1/29) = 3080/24389$ the cutoff in Theorem F.16.3 would become 30. Missing: a written proof of each step of that modified argument; the exact cutoff between 13 and 30 (or 34) is open. [7 October 2026; sources 54–59] *Advanced by source 59* for targets without involutions: the formula holds for $N \geq 16$ (Corollary II.251.3), from source 11's cubic and the published correction theorem, as credited; this supersedes the marked derivation of 30 above. Its $\mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/3\mathbb{Z}$ map (energy $1296 > 1266$) shows the cutoff is at least 13. Source 59 asks (its Question 12.9) for the maxima and extremizers for $2 \leq |G| < 16$, in particular whether the formula holds at 13, 14, 15. Targets with involutions keep the cutoff 34. [7 October 2026; sources 60–69] *Settled for elementary abelian 3-group targets by source 67*: the cutoff is 13, sharp (Corollary II.295.5; the sharpness example is source 59's). Other targets without involutions remain between 13 and 16, targets with involutions at 34.
- II.22** *Erasures against corrupted values* (11 Question 5): a sharp two-parameter envelope for a prescribed proportion of missing inputs and of corrupted remaining values; deletion and corruption have different second-order costs.
- II.23** *Efficient certified recovery* (11 Question 6): recover the canonical affine map in nearly linear time with an a posteriori certificate matching the cubic bound. [6 October 2026; sources 15–20] Source 17 (its Question 9) asks the same for restriction: estimate the signed objective and the exceptional mass without enumerating all tuples, with a runtime theorem. [6 October 2026; sources 28–33] Source 29 (its Question 8) asks the same for its joint theorem: reconstruct Q, ψ, ρ from approximate derivative data with a checkable certificate. [7 October 2026; sources 47–53] *Source 17's question is answered in its ambient-size form by source 48*: exact collision-aware inventories round a signed combination of energies deterministically in $O(C_s n D)$ operations without enumerating tuples, and source 17's restriction is constructed with its retention (Theorems F.53.1, II.194.1); near-linear and $\log D$ -polynomial versions are open (item II.54). [7 October 2026; sources 54–59] Also asked by source 56 (its Question 12.9): a nearly linear or sublinear decoder for relative repair, and an efficient weighted quadruple sampler on an explicit Fourier-uniform domain. [7 October 2026; sources 60–69] Source 67 gives the sharp unique-recovery thresholds $1/4$ and $4/9$ for elementary abelian 3-group targets (Theorem II.295.3) and asks (its Question 7.6) for the optimal query and sample complexity at those thresholds.
- II.24** *Higher-order arrangements: extremal geometry, label repair and stability* (11 Question 7; 14 Questions 5–6). Analogues of the affine-rigidity theory for the multilinear arrangement counts, tracking repeated coordinates and torsion (11); a Hamming-distance repair of labels once the domain is near a product coset (14, first for $\mathbb{F}_p^2$); hypotheses under which near equality in the profile-moment bound yields a common coordinate subgroup, coset-supported section indices and approximately additive top interactions (14). Related to item II.3. [6 October 2026; sources 15–20] Source 15's Questions 2–3 are source 14's Questions 5–6 (label repair at $O_k(\varepsilon)|A|$ points, with coordinate-omitting corrections; robust joint normal forms when $1 - C/\Gamma$ and $R_8 C^6/R_2^7 - 1$ are small). [6 October 2026; sources 28–33] *Advanced on full*

products by source 33: high-agreement repair to $a(x) + B(x, y) + \ell(y)$ with the sharp first and second coefficients $1/s$, $2(s-1)/s^2$ (Theorem F.38.2), the sharp cubic for $s = 4$, erasures (Theorem II.142.1). Its further questions are item II.49; the sparse case is item II.25.

- II.25** *Reaching the high-energy endpoint inside the proof* (11 Question 8): a localization or amplification that reaches respected-quadruple density near one without losing more than it gains, with the retained domain and the normalization tracked together. [6 October 2026; sources 28–33] Source 33 (its Question 6) asks the sparse version of its repair theorem: replace near-full density by quantitative mixing (approximately uniform marginals of conditioned arrangement queries, enough common witnesses), with polynomial dependence instead of the invalid-arrangement error $2s\tau$; the most direct route to Gowers’s Sections 12–13. [7 October 2026; sources 54–59] *Advanced by source 56* in the affine, quadruple-test form: relative repair on any domain with $\|w - \alpha\|_{U^2}^4 \leq \alpha^5/1024$, a condition independent of the error (Theorem F.61.1); source 33’s arrangement version is not addressed. Source 56’s other questions are item II.59. [7 October 2026; sources 60–69] Also asked by source 60 (its Research question 13): mixing assumptions on $A \subseteq G \times H$ under which repair is polynomial in the conditional failure rate, the route from its full-product profiles to Gowers’s sparse steps.
- II.26** *Extremizers and stability at the extension boundary* (12 Questions 1–2): classify the configurations failing extension with $3\varepsilon + 2\kappa = 1$ (are they governed by a three-part quotient as in Theorem II.35.3?), and the structure forced near the boundary under nondegeneracy assumptions. [6 October 2026; sources 28–33] Source 30 (its Question 6) asks how far its budget $(2k-1)\kappa + 2k\eta < 1$ can be improved; at $k = 2$ source 12’s sharp Theorem II.35.2 already improves it (note after source 30’s Theorem II.131.3).
- II.27** *Prime cyclic groups with the same target* (12 Question 3): the optimal deterministic density-overlap region for $G = H = \mathbb{Z}/p\mathbb{Z}$, especially for Bohr sets with a common spectrum; the universal construction does not settle it.
- II.28** *Bohr radii: constants, measured growth, anisotropy* (12 Questions 4–6): the best $c(a)$ in finite radius selection ($c(7/8) = 1/8$ is proved); carrying the measured growth ratio instead of 4^d through later stages; separate radii for different characters weighted by $|\hat{g}(\gamma)|^4$. [6 October 2026; sources 28–33] *Advanced by source 30*: at a fixed radius the translation loss is controlled by the measured growth $|B(3\rho/2)|/|B|$ (Theorem II.129.3), with unequal widths allowed, and exponential dependence on the rank is necessary (Corollary II.135.5). It asks (Questions 4–5) whether the additive structure of the spectra produced in Section 10 forces D_B far below 3^r , and how to choose anisotropic widths before the map is constructed (merged with source 12’s Question 6).
- II.29** *Directional repair* (12 Questions 7–8): hypotheses under which a small global proportion of bad quadruples forces the uniform collision hypothesis of Theorem II.36.2 (directions with few representations are the difficulty; compare source 07’s caveat that collision is not raw energy), and the best repair region in $(\varepsilon, \kappa, \eta)$, including whether $1 - \chi(\chi(\eta))$ is attained.
- II.30** *Measure and infinite-group analogues* (12 Question 12): translation loss and synchronization for finite-measure subsets of locally compact abelian groups.
- II.31** *The range and the examples of the coset envelope* (14 Question 4): the best universal range of ε for $|A\Delta C| \leq \tau(\varepsilon)|A|$ (now $1/100$), the range forcing coordinate splitting, and all exact envelope examples. *Partial answer (intake observation, checked at the write)*: let H be a

subgroup, $K \subseteq H$ a coset of a subgroup L , $k = |K|$, $A = H \setminus K$, $m = |A|$ and $\delta = k/m$. Then $|A \cap (A + d)|$ is m for $d \in L$ and $m - k$ for $d \in H \setminus L$, so $E(A) = km^2 + m(m - k)^2$ and $m^3 - E(A) = mk(m - k) = \delta(1 - \delta)m^3$; as $|A \triangle H| = k = \tau(\delta(1 - \delta))m$ for $\delta \leq 1/2$, every such set with $\delta(1 - \delta) \leq 1/100$ attains the envelope (source 14's Proposition II.43.4 is the case $H = \mathbb{F}_p$, $K = \{0\}$). [6 October 2026; sources 15–20] Source 15 (its Question 1) asks for the best uniform $\delta(\varepsilon)$ with a Cartesian coset within $\delta(\varepsilon)|A|$ when $Q(A) \geq (1 - \varepsilon)|A|^3$: *its leading constant is 1*, by source 14's Theorem II.43.3 and Proposition II.43.4 (= source 15's Example II.52.3); the largest valid range is this item. [7 October 2026; sources 40–46] *Advanced by source 45*: the envelope holds for $\varepsilon < 1/9$ in every abelian group, with the canonical subgroup $\{d : |A \cap (A + d)| > m/2\}$, uniqueness, and all positive equality cases, the punctured cosets $C \setminus K$ with K a subcoset of index at least nine (Theorems II.185.4, II.185.8); a nested-hole envelope with leading constant one (Theorem II.185.11); collision equality forces one punctured coordinate (Theorem II.185.15). Source 45 asks (its Questions 6.6–6.10): the largest range ε_* , and separately the largest ranges for uniqueness and for the autocorrelation threshold (at $\varepsilon = 1/4$ uniqueness fails on $\mathbb{Z}/3\mathbb{Z}$); whether $q(d)(1 - q(d)) \leq \varepsilon$ holds universally (proved with 2ε ; equality on punctured cosets shows that no constant below one works); the maximal range of nested-hole repair and the group-sensitive bound between equality parameters; the best bound for a fixed group and cardinality; stability of collision equality and the labels on near-equality domains. [7 October 2026; sources 60–69] *Answered by source 68* in its range and equality parts: the subgroup and the envelope $\tau(\varepsilon)$ for $\varepsilon \leq 1/4$, optimal even in cyclic groups; uniqueness of the closest coset for $\varepsilon < 2/9$, optimal; equality exactly on punctured cosets of index at least four; $q(d)(1 - q(d)) \leq \varepsilon$ with coefficient one; nested repair with inner defect up to $1/4$ (Theorems II.297.1, II.298.1, II.298.5, II.299.4, Corollary II.298.4). This answers source 45's Questions 6.6–6.8 and source 14's Question 4 in those parts; coordinate splitting and Cartesian structure are not resolved. Source 68 asks (its Questions 8.2–8.6): the discrete envelope $\Phi_{G,m}(u)$ for a fixed group and cardinality (source 45's Question 6.9); stability across the quarter-defect transition; equality in the multilevel inequality; a weighted coset theorem; nonabelian analogues and the equality cases of the mixed remainder.

Added in the fourth write, 6 October 2026 (sources 15–20):

- II.32** *Optimal product-property amplification* (15 Question 5). Are the power $3d - 2$ of χ and the exponents of Theorem II.49.6 optimal under PP_γ alone, or can the reference-cube averaging exploit compatibility between side vectors (starting from Proposition II.53.3)?
- II.33** *Weighted near-extremizers of the cubic cap* (15 Question 7): convert near equality in the weighted cap into a rounding theorem under lower-density or nonconcentration hypotheses; exact saturation forces zero–one weights (Theorem II.56.1 and the remark after it).
- II.34** *Nondegenerate and anisotropic arrangements* (15 Question 8): the error term needed when arrangements must have distinct vertices or side vectors from a prescribed family; a fixed side restriction keeps the moment identity, excluding cross-cube coincidences need not.
- II.35** *Restriction for the higher arrangements* (17 Question 5): the complete short-relation lattice of the configurations of Sections 12 and 15, and a Fejér-type filter controlling every relation class; the missing step before source 17's gain can be propagated. [7 October 2026; sources 60–69] Also asked by source 62 (its Question 12.3): a weighted Smith normal form for the constraint matrix of a higher arrangement.

- II.36** *Spectral data and exactness under restriction* (17 Questions 6–7): keep prescribed Fourier information while restricting, and upgrade a $1 - \eta$ respected proportion to an exact Freiman homomorphism of the same order without losing it. *Partly answered* for exactness from quadruple energy, with polynomial loss, by source 10’s Corollary II.21.2; the spectral part is open. [7 October 2026; sources 60–69] Also asked: a restriction compatible with prescribed large Fourier coefficients of the original domain (62 Question 12.4); exact cells retaining a family of derivative coefficients and the compatibility of cross-sections (61 Question 12.6).
- II.37** *The error-rank distribution* (17 Question 8): inputs where most unrespected tuples have high error rank, combining a small low-rank tail with Theorem II.63.5. [6 October 2026; sources 28–33] Source 29 (its Question 9) asks to optimize its exact signed certificate over the full rank profile $(G_{r,e}, U_{r,e})$.
- II.38** *Degenerate arrangements beyond first order* (18 Questions 1–5). The coefficient of p^{-2} (first $d = 2$, small k), equivalences among nonparity loci, the full moment-order distribution, the intersection lattice of the C_d balanced-subset hyperplanes, small primes $p \leq 2d$ and the case $d = 1$. The placements of sources 22 and 27 record the leading coefficient for $d = 1$ and the second coefficient for $d = 2$; those are for the writes that add them. [6 October 2026; sources 21–27] *Answered for $d = 2$* : the coefficient of q^{-2} is $(k^2 + 21k - 6)/2$ for characteristic $p > 4 \cdot 2^k$ (source 27’s Theorem F.32.3); for general d it is $S_{k,d} - \binom{k}{2} - kL_d - b_{2,d,p}$ with one hyperplane statistic $b_{2,d,p}$ unknown (its Corollary II.99.2). The same question is source 22’s Question 1 and source 25’s Question 1. *The case $d = 1$ is answered in its leading coefficient*, $(k^2 + 3k)/p$, by source 22’s Theorem II.85.1, which also answers the leading-coefficient part of source 27’s Question 10 (the single-pair regime). *Small characteristic*: the first-order count in every characteristic is source 27’s Theorem II.96.3, which answers the counting part of source 25’s Question 4 (the classification of the additional components in characteristics two and three is not given) and advances source 22’s Question 6 (its balanced count for $p \leq 4d$). Still open: $b_{2,d,p}$ for $d \geq 3$; intermediate characteristics $5 \leq p \leq 2k + 2$ (27 Question 4); exact counts in characteristics 2, 3, 5, 7 (27 Question 5); exactness for $d = 1$ (22 Question 2, item II.42). [7 October 2026; sources 40–46] *Source 27’s two questions are answered by source 44*: $b_{2,d,p}$ is characteristic-independent for $p > \max(2d, 16)$ and computed for every d in $O(d^3)$ operations from the nine-point grid, $b_{2,8} = 3,280,234,631,541$ (Theorems II.178.1, II.179.1, II.180.1); the core directions are classified for every $p \geq 5$ (Lemma II.175.1), and the core union is counted exactly in every characteristic (Theorems II.176.2, II.176.3). Still open: third-order terms, exact counts beyond the core in small characteristic, exactness for $d = 1$; source 44’s own questions are item II.52. [7 October 2026; sources 47–53] Exactness for $d = 1$ is answered for $k \leq 3$, with the second coefficient in every dimension, by source 49 (item II.42).
- II.39** *Exact sampling and interpolation* (19 Questions 1–5): classify the maximizers of the finite minimax (how many distinct class sizes); explicit rates in the sharp law $r\mathcal{M}/q \rightarrow C_d$; anchor selection weighing lost mass against slice cost; going below the scale q/σ with algebraic relations among the rows; robust lifting when rows are only approximately polynomial. [6 October 2026; sources 34–39] *Answered for source 39’s rule, open for source 19’s*: for the rule that keeps sampled points the maximizers are balanced and capped (source 39’s Theorem II.155.3, with an explicit cap for affine classes); for source 19’s rule the same formula fails in 199 of 2,835 small cases (note there). Rates are advanced by source 39’s two-regime limit (Theorem II.157.3). Source 39 asks (its Questions 1–6, 8 and 9): every maximizing profile and attaining row family; explicit finite corrections to the golden-ratio law; the first maximizing size for $t > 2$; t growing with the sample; algebra-aware recovery rules with their own bounds; compressing

labelled candidate lists; nonuniform column costs; robust recovery when the good classes are unknown. [7 October 2026; sources 60–69] Source 67: $E_t(L_1, \dots, L_r)$ candidates suffice and are necessary for the interface (Theorems II.339.1, II.339.6; the affine case formalized since 081abc4d3). It asks (its Question 7.7): with $u > t$ valid anchors per class, which compatibility hypothesis improves E_t , and how the optimum depends on u and the list profile.

- II.40** *Proper modular boxes without inherited coordinates* (19 Question 6): replace the inherited-coordinate hypothesis of Corollary II.70.8 by the proper-box predicates of the Lean development, with an explicit width loss. The formal local lift (`section16_local_affine_lift`, `Proofs16LocalAffineLift.lean`, line 13, at 0b4890a1c) avoids the hypothesis with the weaker width $\sqrt{(m/8)^a}/4$ (note after that corollary). [6 October 2026; sources 21–27] *Advanced by source 23*: its Theorems II.103.6 and II.103.7 need no inherited-coordinate hypothesis and keep width $L_m^A \geq (m/3)^A$ on proper modular boxes, conditionally on the slice covers they assume; whether the formal interface accepts their long final axes is not checked (a formalization candidate). Source 26 (its Question 6) asks how its integer-box almost-cover (Theorem I.44.1) can meet the exact-partition and base-loss requirements of this interface.
- II.41** *A consistent family of structural controls for Lemma 16.10* (19 Question 10): controls that keep the outer parameter dependence, absorb the shared-slice output and are stable under restriction, union, sums, an unused coordinate, slices and interpolation; both displayed closing comparisons fail (Section V.80.2; formally `section16_same_dimension_closing_comparison_fails`). `lemma_16_10` is open at HEAD. [6 October 2026; sources 21–27] Also asked by source 23 (its Question 11: substitute the actual class, anchor and budget counts into its gluing theorem and find control functions closed under every step, with a checked recurrence). [6 October 2026; sources 28–33] Also asked by source 28 (its Question 13: an altered induction invariant that closes the Section 16 comparison with synchronized proper partitions). [6 October 2026; sources 34–39] Also asked by source 39 (its Question 10: replace the invalid closing comparison by a valid dimension-dependent recurrence for scale, graph count and exceptional mass, with endpoint-safe controls carrying a slack factor above one). `lemma_16_10` is still open at HEAD 89c1cbc38. [7 October 2026; sources 47–53] *Answered negatively for the catalogue form by source 50*: the two-premise `lemma_16_10` is itself false, not only its closing comparison (Theorem II.233.1); a repair must strengthen the input (the product property, whose full-domain endpoint is multiaffinity, Proposition II.234.1) or change the output controls (source 49’s acyclic recurrences, Theorem II.235.5). The repository has formalized the negation and keeps `lemma_16_10` open (Section V.1). Also asked by source 50 (its Research questions 1–2: for $0 < \gamma < 1$, the quantitative information the relation product property retains after selections and restrictions, and whether it controls one-coordinate summands at the needed branch budget; a closed recurrence for graph counts, width exponents and iteration parameters from the explicit local lift, with the least extra dependence on the input branch count, uniform words requiring $(1 - \eta - o(1))R$ graphs) and source 49 (its Question 10.14: replace I7 by a sharper simultaneous multilinear recurrence, and find which improvement changes the final width exponent).

Added in the fifth write, 6 October 2026 (sources 21–27):

- II.42** *The structure of the degeneracy locus* (22 Questions 2–5; 25 Questions 2, 3, 7; 27 Questions 1–3, 6, 7). Is $\delta_{k,1}(p) = L_k(p)$ for all large p (22; the intake found equality for $k \leq 3$ at the primes it tested, evidence only)? Degree-two Boolean certificates and a feasible-class filtration in every degree with a valid summation by parts, and a generating function for $J_{k,d}$ (22); a much

smaller second-order constant than $U_{k,d}$ (25; answered in substance by source 27's $S_{k,d}$ with a cubic remainder, Theorem II.97.8); the intersection lattice of the cross-section hyperplanes, which is source 27's statistic $b_{2,d,p}$, first for $d = 3$ with its 51 hyperplanes (25, 27); arbitrary defining relations w (25); a third-order structural theorem and the codimension in terms of the minimal Boolean degree (27); whether the count stays polynomial beyond eight vertices (27); minimal moment subfamilies keeping the asymptotics (27). [7 October 2026; sources 47–53] Answered for $k \leq 3$ by source 49: $\delta_{k,1}(p) = L_k(p)$ at every odd prime for $k \leq 2$ and, by a finite certificate, $k = 3$ (Corollary II.222.5, Theorem II.223.1); in every dimension the second coefficient (Theorem II.222.4), an exact count in characteristic three and eventual Laurent polynomiality. Source 49 asks (its Questions 10.1–10.7): exactness for $k \geq 4$, first $k = 4$; the exact coefficient of p^{-3} ; a classification of derivative rank three; a sharper finite remainder than $M_k^2/2$; the characteristic polynomial χ_4 and the bad-prime sets; a short proof of the dimension-three certificate explaining the primes three and seven; the multiple-slice second coefficient for $d \geq 2$ with both moment constraints.

- II.43** *Efficient degeneracy testing* (25 Question 5; 27 Question 8). Detect a nontrivial ternary moment relation without scanning 3^n patterns or 2^m subset sums, using the $k + H_d$ hyperplanes and the core family as preprocessing, and return a short verifiable certificate of degeneracy or nondegeneracy.
- II.44** *The selection step of Lemma 15.5* (25 Question 6; 27 Question 9). A nonnegative selection kernel distinguishing respected arrangements by more than the ratio $1 + 2^{1-n}$ while keeping manageable degeneracy conditions (25; the exponential loss of (II.92.12) is not removed by better counting); bounds by the actual cross-section distribution of a dense collection instead of the ambient union bound, avoiding the loss of Proposition II.101.1 (27). [6 October 2026; sources 34–39] Source 35 (its Question 4) asks which nonnegative trigonometric polynomial $K \leq 1$ with nonnegative coefficients best balances the universal signal, the coefficient-stratum mass and the point marginal; source 40's placement records a kernel-class barrier, for its write. Source 35's degree-weighted selection itself is in Section II.144; its other questions are item II.50. [7 October 2026; sources 40–46] Source 40 is now printed (Section II.164). It asks (its Question 1) for the optimal finite-bandwidth kernel for fixed s , H , Q and how close Fejér is, with certified positivity; source 44 (its Question 6) for an analogue of its core product for arrangements sampled from a dense set, with explicit conditional-concentration hypotheses.
- II.45** *Common base partitions at Lemma 16.10* (23 Questions 9–10). Avoid the product of width exponents over all sampled slices in Theorem II.103.6 by a joint structural theorem giving one common partition directly; and a width-dependent localization rule that improves the constant $1/3$ when a large lower bound for the child width is known. [7 October 2026; sources 47–53] Also asked by source 49 (its Question 10.13: model all anchor slices together with a width loss polynomial, not exponential, in the number of anchors, at the factor c_H^r of (II.235.28)) and in the question paragraph closing its Section 8 (Section II.235.11: simultaneous slice refinement, a smaller minimum-width threshold keeping the global core independent of the local tolerance, a sharper I7, and recurrences as expression trees for a formal implementation).

Added in the sixth write, 6 October 2026 (sources 28–33):

- II.46** *The critical-norm profile beyond the sharp range* (31 Questions 1–4 and 12). The optimal universal relation between $\kappa(f)$ and the best coset-phase fidelity beyond $1 - \kappa < 1/120$, and where other configurations take over from the two-state curve; the optimal odd-order coefficients (Theorem II.107.3 excludes the order-two mechanism; source 38's placement

records the odd-order modulus, for its write); a quantitative classification of near-equality states of the quartic envelope; the exact feasible region of $(\varepsilon, t, z, \beta)$ in Theorem II.108.1; equality in the mixed derivative inequality for two functions, and analogues for higher Gowers norms. [6 October 2026; sources 34–39] *The optimal odd-order coefficients are source 38's:* $\mathfrak{S}_{\text{odd}}(\varepsilon) = \varepsilon/4 + 3\varepsilon^2/16 + (\sqrt{2}/64)\varepsilon^{5/2} + O(\varepsilon^3)$, and without the fractional term when $\gcd(|G|, 6) = 1$ (Theorem II.150.11). Source 38 asks next (its Questions 7–10): the ε^3 coefficient and whether the three-point family stays extremal; the coprime-to-six cubic term; a constructive reconstruction from sampled derivative data with near-sharp constants; sharp critical stability for U^d , with critical exponent $2^d/(d+1)$. [7 October 2026; sources 54–59] *Extended by source 55* to bounded functions and every degree: squared distance to phase polynomials $(2/q)\varepsilon + b_d\varepsilon^2 + O_d(\varepsilon^3)$ with $b_d = (3^d - 1)/(2q^2)$ (odd order) and $(q^2 - 1)/(3q^2)$ (even order), both sharp; at $d = 3$ the coefficients 13/64 and 21/64 of sources 38 and 31 (Theorem II.237.1). Source 55 asks (its Questions 12.1–12.3): the optimal ε^3 coefficient and whether finitely many torsion statistics determine it; whether the one-character even-order curve is exact for $d \geq 3$; quantitative rates in its tangent rigidity theorem.

- II.47** *Logarithmic selection at Section 10* (28 Questions 1–6). Are both logarithmic losses in $\alpha p(\alpha)/20$ necessary ($\Omega(\alpha/\log(1/\alpha))$ or $\Omega(\alpha)$ with a comparable original-domain spectrum)? Can $O(\alpha^{-3})$ frequencies be reduced, or counted by an independent generating set? The best simultaneous dependence on the defect, the retained fraction and the radius; output order eight or any fixed order; which conclusions survive input order four or two; a certified finite construction faster than enumeration. *Marked question (the intake's observation, in no source; not checked at the write):* does Theorem F.33.1 with Corollary II.126.2 give the conclusion of the formal `theorem_10_13` (Section10.lean, line 473) with the formal spectrum `domainLargeSpectrum` at λMN , a larger retained set ($1.08 \cdot 10^{-4}$ against $1.07 \cdot 10^{-9}$ at $\alpha = 1/6$) and a larger radius? Missing: a check that the order-eight defect 2^{-21} of source 28 is measured as in `DomainApproxHomOfOrder`, and that $\text{Bohr}(\Gamma, \zeta_\alpha)$ contains the formal Bohr set at `section10Zeta`.
- II.48** *Fixed-radius Bohr invariance* (30 Questions 1–3, 7, 10–11). The exponential rate of $\mathcal{L}_r$ ($\sqrt{3}$, 3 or another base); the prime cyclic case, where the Chinese-remainder obstruction is unavailable; joint sharpness of escape and packing; whether the 5/14 compatibility threshold is optimal and whether target-fibre information improves $\theta < 1/2$; constraints beyond rectangles (convex constraints, covering numbers in place of 3^r); efficient exact certificates for D_B below 3^r .
- II.49** *High-agreement repair of arrangements* (33 Questions 1–5 and 7–10). The optimal common-prime quadratic coefficient, between $(3s-2)/(4s^2)$ and $2(s-1)/s^2$ (23/512 and 15/128 at $s = 16$); the optimal cubic coefficient for $s \geq 6$; equality and near-equality in the sharp cubic; the activation threshold $1/(240(s-1))$ and the global feasible pairs (δ, ε_4) ; sharp bounds from row profiles; nonzero increments and nonrepeated vertices; two-direction testing with repair to a biaffine map; higher multilinear arrangements; fast decoding with certified sample complexity. [7 October 2026; sources 60–69] *Answered by source 60:* Question 2 (every coefficient of the repair bound is optimal in succession), Question 3 in its equality part ($a\mathbf{1}_C(x)\lambda(y)$, C of index $n \geq 3$, at $\delta = 1/(2n)$) and Question 8 locally (the modulus $J(\eta)$ and biaffine repair through degree $s-2$); Question 1 advanced to $[(s-1)/s^2, (3s-2)/(2s^2)]$ (Section II.331). Source 60 asks (its Research questions 5–9): quantitative near-equality; the common-prime coefficient in the new interval; the finite-error envelope and the activation threshold; the optimal two-direction coefficient at degree $s-1$ (the cubic at $s = 4$) and equality in the intersection theorem; three or more coordinate directions, beginning with triaffine models.

Added in the seventh write, 6 October 2026 (sources 34–39):

- II.50** *Degree-weighted arrangement selection* (35 Questions 2, 3, 5, 6 and 14). Refine the coefficient strata by the rank of the matrix $(A_{j,T})$ so that several low Boolean moments restrict simultaneously; the exact mass of the balance spaces against μ^{s-D_t} ; fixed characteristic with $\alpha\eta \rightarrow 0$, and targets with small torsion; an intrinsic lower bound on the field size for any selection with the stated purity and count; an efficient algorithm finding B' . Section II.144; the kernel question is merged into item II.44. *Marked question (the intake's observation, in no source; not checked at the write):* does Corollary II.145.3 give the conclusion of the formal lemma_15_5 (Sections14_15.lean, line 414) with an explicit N_0 and retention $\alpha^s \eta^{s-1} / (2(s+2)^s)$, $s = 2^{k+4}$? Missing: a check that the ordered and degenerate conventions of `respectedGeneralArrangementCount` and `generalArrangementCount` match source 35's arrangement count, and that its purity $1/(1+\eta)$ gives the formal $(1-\eta)$ form.

Added in the eighth write, 7 October 2026 (sources 40–46):

- II.51** *One-stage Fejér selection* (40 Questions 2–5). Improve the exceptional constant $C_{k,s,H}$ by identifying scalar-equivalent or coincident zero sets of the coefficient vectors and their intersections (zero sidelengths and coincident vertices kept); whether a field size of order $\beta^{-(2d-1)}(s/(\alpha\eta))^{2s}$ is necessary for a one-stage theorem of this kind or can be reduced at the same mass, for example by stratifying extra-relation ranks (source 35's degree weighting is one such reduction); comparable selection in characteristic $p \leq 2(L-1)$, separating coefficient collisions from target torsion; configurations whose feature relations form a lattice of rank greater than one. Section II.164; the kernel question is merged into item II.44 and the class barrier into item II.7.
- II.52** *Beyond the second Whitney coefficient* (44 Questions 1–5 and 7). The third Whitney coefficient from a fixed grid in $\{-1, 0, 1\}^3$ with a finite rank-three incidence weight, in polynomially many operations; which non-core directions contribute at order q^{-3} on the regular region; short certified recurrences for the eight lattice energies and their combination B_d ; which primes actually change $b_{2,d,p}$ and whether the cutoff 16 is optimal (the intake found the rational value already for $p \geq 7$ at $d = 3$ and $p \geq 11$ at $d = 4$); the largest k for which nondegenerate d -arrangements exist over a fixed field, against the necessary conditions of Corollary II.176.4; the coefficient alphabet $\{-M, \dots, M\}$ in place of $\{-1, 0, 1\}$. Section II.173; source 44's question on relative distributions is merged into item II.44.
- II.53** *The top of the torsion-free energy spectrum* (45 Questions 6.1–6.5). Every set at the third level $M_m - 8(m-3)$ for large m (two models are known: $\{0, \dots, m\} \setminus \{2\}$ and $\{0, 2, 3, \dots, m-1, m+1\}$) and the smallest valid threshold; positive real sets with bounded Schur defect s ; sharp edit or containment bounds when $D(A) = O(m)$ or $o(m^2)$, for example whether $D(A) \leq Km$ puts all but $O_K(1)$ points into one progression of length $m + O_K(1)$; a weighted ordered-completion identity; effective hypotheses on $A \subseteq \mathbb{Z}/p\mathbb{Z}$ that allow the three levels to be transferred by a Freiman 2-isomorphism. Section II.183. Source 57, placed by 7dc1ee29d and not yet written in, records the complete third level (J_m and one further class for $m \geq 7$) and bounded Schur deficits (for its write). [7 October 2026; sources 54–59] *The first two questions are answered by source 57:* the third level is $\{J_m, K_m\}$ up to affine maps for $m \geq 7$, with one more class at $m = 6$, and bounded Schur deficits are integer partitions for $n \geq 5s - 1$ (Theorems II.292.7, II.292.3, II.292.4). Source 57 asks (its Research questions 1–5): the optimal Schur threshold $N(s)$ (known $N(1) = 4$, $N(2) = 6$) and the exceptions below it; the fourth and later levels with their thresholds; energy formulas for partition models; whether $D(A) \leq Km$ forces all

but $O_K(1)$ points into a progression of length $m + O_K(1)$; rectification to cyclic groups and weighted versions. [7 October 2026; sources 60–69] Advanced by source 66: bounded Schur defects are partitions for $n \geq 3r + 1$, so $2r + 2 \leq N(r) \leq 3r + 1$ (with source 57's $N(1) = 4$, $N(2) = 6$); the exact hole energy law and the endpoint envelope $M_m - 4r(m - 1 - r)$; the third level from $m \geq 8$ (source 57: $m \geq 7$) and the fourth, $M_m - 8m + 20$ with two templates, from $m \geq 11$ (Theorems II.302.1, II.304.1, II.305.2, II.308.2); this answers source 57's Research question 2 at the fourth level and 3 for hole models. Source 66 asks (its Conjectures 13.1, 13.3 and Questions 13.2, 13.4–13.9): $N(r) = 2r + 2$; the critical-size exceptions at $n = 2r + 1$; the second boundary-energy constant $r^2 - 2r + 4$; the weighted partition spectrum $F(z, w)$; the fifth and later levels; an inverse theorem at linear total defect (source 57's Research question 4); robustness to perturbed relations; weighted and multiset versions; rectification of near-extremal cyclic graphs (source 57's Research question 5).

Added in the ninth write, 7 October 2026 (sources 47–53):

- II.54** *Constructive restriction beyond ambient size* (48 Research questions 1–8). A translated Fejér parameter triple with positive surplus in $\tilde{O}_s(np^2)$ or $\tilde{O}_s(p^3)$ time with the guarantee of Theorem II.194.1 (the total cost is $O_s(np^5)$); subquadratic storage for graph data, with a support-growth theorem; polynomial dependence on list size and bit length for sparse or high-dimensional data, and on the solution geometry for general systems; prescribed size after a heterogeneous Bernoulli filter, with a degree- $2s$ support state; surplus retained together with prescribed Fourier correlations or spectral energy, with a valid separating certificate; a joint filter–precision–search frontier at a fixed computational budget; a faster exact array update with a bit-cost bound; the support profiles realizable by graph energies, and threshold-free restriction from them. Section II.186. [7 October 2026; sources 54–59] Also asked by source 58 (its Question 13.8): a qualifying subset and an a posteriori certificate in time well below enumerating all b^m tuples. [7 October 2026; sources 60–69] Also asked: conditional expectations in time controlled by the sparse input (62 Question 12.7); separation certificates and large exact cells from compressed input (61 Question 12.8); a qualifying subset without enumerating all b^m tuples, and which later steps dominate (68 Question 8.13).
- II.55** *Section 13 interfaces* (51 items 1–4; 52, three paragraphs; 53 items 1–3). Export the actual Stage 13.6 length $\ell R \in \{m, m + 1\}$ (or $\ell R \leq m + 1$) through the wrappers and record short-carrier rectification with reversal and the balanced aligned-parent partition (51 items 1–2, 53 item 1; at HEAD 06b30ca52 the wrapper exports $m \leq \ell R \leq \ell Q$, not the upper bound); reuse the all-scale coefficient localization at unchanged density and state the direct Stage 13.8-to-square result with its effective width (51 items 3–4); a smaller threshold from only the width and short-parent budgets of the sharper route (53 item 2); the best explicit threshold for side $N^{\lambda\gamma(a)}$, $\lambda < 4$ (53 item 3); at Stage 13.7 over composite moduli, the combined package of a unit common step, genuine initial stages, $\ell R \leq \ell Q$ and all six budgets; whether coordinate-affine conclusions are compatible with the later stages, ambient coefficient extraction, translations and bilinear conclusions; a slope in a module or a compatible homomorphism on a difference subgroup with explicit extension data, tested by the two-point congruence criterion (52). Section II.195; 53's item 4 is in item V.1. [7 October 2026; sources 54–59] Source 57 asks (its Research questions 7–8) for a better prepared frequency graph than density $a^{2,996,201,429}$ (a density $a^{2^{30}}$ would lower the final exponent) and for a closed bound on the inherited threshold $N_{\text{sq}}(\delta)$; the second is answered by source 53's threshold, and the repository has its own explicit formula (note after Remark II.290.4).

II.56 *Branch-cover profiles* (50 Research questions 3–9). Interfaces carrying the whole profile $q \mapsto M_q$ instead of one branch count, with composition rules under union, interpolation and common refinement; the shortest universal balance scale on modular progressions, for fixed and for growing alphabets; the true transition to constant optimal lists of degree d (relative balance against nonconstant-graph avoidance); the list-cover optimum when $b_{(q)} < dR$, by a computable incidence invariant; near-linear certification of balanced words, implicit polynomial avoidance, or an explicit algebraic word balanced at logarithmic scale; composite rings and fibre multiplicities; stability and reconstruction from noisy or partial observations. Section II.226; 50's Questions 1–2 are in item II.41 and its Question 10 in item V.2.

Added in the tenth write, 7 October 2026 (sources 54–59):

II.57 *Common phases at Proposition 8.1* (54 Questions 11.1–11.9). The best balanced-indicator coefficient in place of $1/(B\sigma)$, and whether the balanced constraint beats the amplitude-only bound as $N \rightarrow \infty$; the exact stability modulus $\Phi_m(\varepsilon)$ for cycles of order $m \geq 5$; repair preserving a pointwise cap; sharp ceilings localized to an interval or Bohr set; one frequency good for a whole family of windows; a common global phase in higher-degree analogues of Proposition 8.1; nonsymmetric selectors on a restricted shift set; stability near the clique ceiling for Cayley graphs Γ_S ; a faster certified selector than the $O(N^2 \log N)$ scan. Section II.267; 54's Questions 11.10–11.11 are in items V.1 and V.2.

II.58 *The polynomial-phase endpoint beyond finite groups and exact sampling* (55 Questions 12.5–12.9). The sharp expansion for the interval norm $U^d([N])$ with boundary terms; joint distance for mixed cubical near-extremizers, with compatible phase tuples; sampling laws with approximately independent triples; second-countable compact abelian groups, with parity replaced by an order-two character; the Pauli uniformity norms of unitary matrices and which part of the two-torsion dichotomy survives. Section II.236; 55's Questions 12.1–12.4 are in items II.46 and II.2, its Question 12.10 in item V.2.

II.59 *Relative affine repair* (56 Questions 12.1–12.6 and 12.8). The weakest U^2 hypothesis $u^4 \ll \alpha^\theta$, $\theta < 5$, forcing error-independent repair; the best coefficient $C(\lambda)$ at fixed uniformity; a torsion-sensitive replacement for the coefficient 12 of t^2 ; the largest repair range beyond $\lambda \leq 1/32$, $\varepsilon \leq 1/1024$; a lower rank than $4r + O(1)$ on quadratic level sets; higher-degree repair with a fixed domain condition; regular Bohr sets or long progressions as domains. Section II.248; 56's Questions 12.7, 12.9 and 12.10 are in items V.1, II.23 and V.2.

II.60 *Character separation and exact restriction* (59 Questions 12.1–12.5). The optimal radius $R_p(q)$ at the exact prime-primary rank (at least $1/(2p)$; $R_3(3) \leq 2/9$); joint optimization of radii and character-image costs; whether $\lambda_k(q) \geq c_k/(q+1)$ for the retained proportion from q sign orbits of defects; a linear exponent $a(k) = O(k)$ in the doubling restriction (now $(2k+1)\log_2(5k)$); and, beside the lonely runner conjecture, whether small doubling $|\Gamma - \Gamma| \leq q^{1/2}|\Gamma|$ forces one character to separate q infinite-order defects at radius $1/(q+1)$. Section II.261; 59's Questions 12.6–12.10 are in items II.19 and II.21. [7 October 2026; sources 60–69] *Source 61* finds the exponent $(2m+1)\log_2(5m)$ independently, partitions the whole domain into exact cells and proves $\lfloor \log_2 m \rfloor + 1 \leq b_m \leq \log_2(5m)$ with p -primary obstructions (Section II.310). *Source 61* asks (its Questions 12.1–12.5, 12.7 and 12.9): $b_m = \lfloor \log_2 m \rfloor + 1$?; the growth of the graph-extraction exponent a_m^* ; an effective defect parameter below K^{2m+1} ; the optimal relation-energy powers (now $O(m^2 \log m)$); exact against approximate respect on composite cyclic groups; cyclic against arbitrary groups; stability of the extremal torsion examples.

The eight items carry the nine questions: 02 Subsections 11.1–11.3, 03 Research questions 3–4, and 04 Q4–Q7. [6 October 2026; sources 06–10] With sources 06–10 the eighteen items carry nineteen questions and one question paragraph: also 07 Questions 1–7, 10 Questions 5–7 and the paragraph at the end of 10’s Section II.21. [6 October 2026; sources 11–14] With sources 11–14 the thirty-one items carry forty questions and one question paragraph: also 11 Questions 1–8, 12 Questions 1–9 and 12, and 14 Questions 4–6; item II.21 also records the intake’s marked question. [6 October 2026; sources 15–20] With sources 15–20 the forty-one items carry sixty-eight questions and one question paragraph: also 15 Questions 1–3 and 5–8, 17 Questions 1–9, 18 Questions 1–5, and 19 Questions 1–6 and 10; items II.6 and II.8 also record marked questions of the intake. [6 October 2026; sources 21–27] With sources 21–27 the forty-five items carry one hundred and four questions and one question paragraph: also 22 Questions 1–6, 23 Questions 9–11, 24 Questions 1–9, 25 Questions 1–7, 26 Question 6 and 27 Questions 1–10. [6 October 2026; sources 28–33] With sources 28–33 the forty-nine items carry one hundred and fifty questions and one question paragraph: also 28 Questions 1–6 and 13, 29 Questions 5–12, 30 Questions 1–7, 10 and 11, 31 Questions 1–12 and 33 Questions 1–10. [6 October 2026; sources 34–39] With sources 34–39 the fifty items carry one hundred and sixty-nine questions and one question paragraph: also 35 Questions 2–6 and 14, 38 Questions 7–10, and 39 Questions 1–6 and 8–10. [7 October 2026; sources 40–46] With sources 40–46 the fifty-three items carry one hundred and ninety-two questions and one question paragraph: also 40 Questions 1–5 and 7, 44 Questions 1–7, and 45 Questions 6.1–6.10. [7 October 2026; sources 47–53] With sources 47–53 the fifty-six items carry two hundred and twenty-eight questions and two question paragraphs: also 48 Research questions 1–8, 49 Questions 10.1–10.7, 10.13–10.14 and the paragraph closing its Section 8, 50 Research questions 1–9, 51 items 1–4, three paragraphs of 52 and 53 items 1–3. [7 October 2026; sources 54–59] With sources 54–59 the sixty items carry two hundred and eighty-one questions and two question paragraphs: also 54 Questions 11.1–11.9, 55 Questions 12.1–12.9, 56 Questions 12.1–12.6, 12.8 and 12.9, 57 Research questions 1–5 and 7–10, 58 Questions 13.1–13.8 and 59 Questions 12.1–12.10. [7 October 2026; sources 60–69] With sources 60–69 the sixty items carry three hundred and thirty-two questions and two question paragraphs: also 60 Research questions 5–13, 61 Questions 12.1–12.9, 62 Questions 12.1–12.7, 66 Conjectures 13.1 and 13.3 and Questions 13.2 and 13.4–13.9, 67 Questions 7.3–7.7 and 68 Questions 8.2–8.13.

Theme III: cube and progression counts (8 questions)

III.1 *The odd-order constant and the higher-dimensional norm gap* (03 Research question 1; 04 Q3). $c_d^{\text{odd}} = \sup \|f\|_{U^2}^4 / \|f\|_{U^d}^4$ over real centered f on odd-order groups (03), equivalently the best comparison of $\|f\|_{U^d}^{2^d}$ with $\|f\|_{U^2}^{2^d}$ for given group orders (04). *Settled for $d = 3$* (dated note after Theorem III.3.4): $c_3^{\text{odd}} = 2^{-1/2}$. Open for $d \geq 4$: monotonicity gives $\|f\|_{U^2}^4 \leq 2^{-1/2} \|f\|_{U^d}^4$ for $d \geq 3$ (04, after Theorem III.3.6), but its sharpness for $d \geq 4$ is not established. The best odd-order leading coefficient is $R_d c_d^{\text{odd}}$. [6 October 2026; sources 06–10] Source 08’s Question V.34.1 asks for the same constant (κ_d^{odd}) and whether $\kappa_3^{\text{odd}} = 1/\sqrt{2}$: yes, by source 04. Source 09’s Research question on the constrained envelope (item III.17) contains the odd-order leading coefficient relative to $\|f\|_{U^d}^4$ as a special case. [6 October 2026; sources 11–14] Source 14 (its Question 11) asks the same for $d \geq 4$: the first surviving centered correlations of U^d , the sharp real mean-zero comparison with lower norms, and torsion-dependent extremizers. [6 October 2026; sources 15–20] Source 18 (its Question 7) asks for the best fourth-order coefficient without order-two characters, “especially in dimension four”; it is settled at $d = 3$ by source 04’s Theorem III.3.4, which source 18 does not cite. [6 October 2026; sources 21–27] Source 21 (its Question 7) asks for the sharp leading real-centered coefficient in $Q_d(\delta + f)$ for $d \geq 4$ and

its first torsion-dependent correction. [6 October 2026; sources 28–33] *Advanced, not settled, by sources 28 and 29* (Section III.66): $c_4^{\text{odd}} \leq (8/35)^{1/4} = 0.6914416$ and $c_d^{\text{odd}} \leq B_d \downarrow 1/2$ (source 29’s Corollary III.69.1; source 28’s weaker $\sqrt{160/323}$), $c_d^{\text{odd}} \geq 1/3$ for every d , so the limit lies in $[1/3, 1/2]$ (source 29’s Proposition III.69.2, answering source 28’s question on a positive lower bound uniform in d); $c_4^{\text{odd}} \geq 0.5587294$ (source 28) and ≥ 0.5587725 (the intake’s exact certificate on $\mathbb{Z}/7\mathbb{Z}$, a review result; numerical searches suggest about 0.55878, evidence only). Real parts of characters are not extremal for $d \geq 4$ (source 28’s Theorem III.70.9), which answers source 03’s sub-question and source 29’s Question 1 in its second sentence. Among sources 01–33 the bracket is $0.5587725 \leq c_4^{\text{odd}} \leq 0.6914416$; sources 34, 35 and 45, placed and not yet written, record upper coefficients 81/16, 5 and 4795/832 (for their writes). Further asked: the exact c_4^{odd} and whether a bounded number of harmonics suffices, the dependence on character orders and attainment on a fixed group, the asymptotics of c_d^{odd} , and stability near true extremizers (28 Questions 7–11); the limit c_∞^{odd} and whether 1/2 is approached, the sharp profile at fixed self-inverse mass, and stability for actual norm near-minimizers (29 Questions 1–4). [6 October 2026; sources 34–39] *Advanced by sources 34 and 35*: $\|f\|_{U^2}^4 \leq (2/3) \|f\|_{U^4}^4$ (source 34’s Theorem III.71.1; coefficient 81/16 in source 29’s normalization) and the weaker coefficient 5 (source 35’s Theorem III.72.1, the same day); for $5 \leq d \leq 8$ slightly better bounds than B_d (source 34’s Remark III.71.10). Among sources 01–39 the bracket is $0.5587725 \leq c_4^{\text{odd}} \leq 2/3$; source 45, placed and not yet written, records $(832/4795)^{1/4} = 0.6454070$. Further asked: the exact c_4^{odd} , its attainment, and the extremizing supports and phases; stronger coupled certificates; the asymptotics and stability for $d \geq 4$ (34 Questions 7–9); whether the full statistic $J(\lambda)$ determines c_4^{odd} , whether an extremizer lives on a finite cyclic quotient, and higher-order gains from the harmonics of D_d of a cosine (35 Q1); the exact c_4^{odd} and bounded Fourier support, the structure near the minimum, and higher-order gains beyond central-binomial moments (35 Q7–Q9). [7 October 2026; sources 40–46] *Advanced by source 45*: $\|f\|_{U^4}^{16} \geq (4795/832) \|f\|_{U^2}^{16}$ (Theorem III.90.1), so $0.5587725 \leq c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$ among sources 01–46, the lower end being the batch-121 certificate; source 45’s own lower example gives 0.5574336; source 49, placed by **f8bc5e2ec** and not yet written in, records 29/4, that is $c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ (for its write). Source 45 asks (its Questions 6.11–6.14): the exact c_4^{odd} , its extremizers and whether it is attained on a finite (cyclic) group; a stronger, possibly sum-of-squares, moment certificate; the Fourier support of an extremizer, whether two character pairs suffice and whether order five is intrinsic; stability, and the constants c_d^{odd} for $d \geq 5$ with possible strict gaps. [7 October 2026; sources 47–53] *Advanced by source 49*: $\|f\|_{U^4}^{16} \geq (29/4) \|f\|_{U^2}^{16}$ (15/2 when a largest coefficient has order other than five) and $\|f\|_{U^5}^4 \geq (50/29) \|f\|_{U^2}^4$ (Theorem III.93.1), so $0.5587725 \leq c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ among sources 01–53 and $c_5^{\text{odd}} \leq 29/50$. Source 49 asks (its Questions 10.8–10.12): c_4^{odd} , or a much shorter interval, and whether the supremum is a limit on cyclic groups with boundedly many Fourier pairs; whether five-torsion changes the optimal constant, and raising the order-five dual profile to the non-five level; optimized dual functions $\phi_1 + t_3\phi_3 + \cdots + t_{2s+1}\phi_{2s+1}$ with exact certificates; envelopes from the two or three largest pairs; direct dual tests for U^d , $d \geq 5$, and the behaviour of c_d^{odd} and near-extremizers as d grows. [7 October 2026; sources 60–69] *The limit is settled by source 63*: $\|f\|_{U^2}^4 \leq (1/3 - 4/(3N^2) + 1/N^3) \|f\|_\infty^4$ on every group of odd order N , sharp on $\mathbb{Z}/N\mathbb{Z}$, and $c_d^{\text{odd}} \leq 1/3 + 10 \cdot 2^{-\lfloor d/8 \rfloor}$ for $d \geq 80$, so $\lim_d c_d^{\text{odd}} = 1/3$ (Theorems III.96.4, III.97.1); source 60, independently, identified the limit with $\beta_{\text{odd}} \in [1/3, 0.3464533951]$ and proved the connected limit 1/3 (Section III.94). This answers source 29’s Question on c_∞^{odd} , source 28’s on the asymptotic constants, and source 60’s Research questions 1, 2 and the first part of 3. At fixed d the bracket is unchanged: $0.5587725 \leq c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$. Asked: $\beta(q)$ on

groups of fixed odd exponent, $\beta(3) \in [56/243, 5/18]$, and extremizers (60 Research question 3, second part; 63 Research question 7.4); stability at the connected constant (60 Research question 4) and for the bounded inequality (63 Research question 7.3); c_d^{odd} at a fixed $d \geq 4$ and its attainment (63 Research question 7.1); the optimal rate (63 Research question 7.2); other reality and centering constraints (63 Research question 7.5).

III.2 *The fifth and sixth centered coefficients* (03 Research question 2). Classify the supports of sizes 5 and 6 for which (III.2.27) has solutions with no trivial character, with multiplicities and the quotient correlations they need; Lemma III.3.3 does this for $d = 3$ on odd-order groups, and (III.2.29) shows that order three enters. [6 October 2026; sources 06–10] *The fifth coefficient is answered by source 09* (Theorems III.20.1, III.21.1, Proposition III.20.2): on every finite abelian group and for every $d \geq 3$, only the star and ternary supports survive, with multiplicities Q_d , S_d and correlations M_5 , $T_{3,5}$. Open: the sixth coefficient, asked also by source 08 (its third question, fifth and sixth orders for general d with their torsion dependence) and source 09 (its Research question 1, a dimension-uniform sextic expansion with a seventh-order remainder); source 09’s Theorem III.23.1 shows that an isolated order-five term must appear there. [6 October 2026; sources 11–14] Source 13 (its Question 10) asks how far the low-order terms of higher cube expansions can be expressed through lower uniformity norms and a few spectral moments, first classifying five- and six-circuits for $d = 4$. [6 October 2026; sources 15–20] The sixth order is also asked by source 16 (its Question 10: primitive relation modules for six Boolean vertices in arbitrary dimension) and source 18 (its Question 6: six-vertex small-matrix classification with invariant factors and a remainder bound). [6 October 2026; sources 34–39] Also asked by source 34 (its Question 6: classify the full-support lattice types at the first unresolved support degrees with their multiplicities as functions of d , and express the sixth and seventh coefficients in a finite basis of Fourier statistics).

III.3 *The secondary extremal coefficient* (04 Q1). Between $(9/4)^{1/3}$ and 6 (display (Q1)); does the limit exist, and is the two-frequency value optimal? [6 October 2026; sources 11–14] *Settled by source 14*: the limit exists and equals $(9/4)^{1/3}$, uniformly over groups of order coprime to 6 (Theorem III.34.1; in source 04’s normalization, Remark W4). The two-frequency value is optimal; the 6 of Theorem III.3.8 stays an explicit upper bound. What remains is in items III.4 and III.23.

III.4 *Stability for cube extremizers* (04 Q2), beyond the U^2 stability of Corollary III.3.5, and whether near-optimal secondary behaviour concentrates on $\{\pm\xi, \pm2\xi\}$. [6 October 2026; sources 11–14] *Answered for the second part*: near-optimal secondary behaviour concentrates on $\{\pm\xi, \pm2\xi\}$ in U^2 , with the amplitude and phase forced (source 14’s Theorem III.34.3), and not in L^2 (the remark after it). Source 14 (its Question 3) asks for an effective version, with explicit dependence on the deficit and u/δ , and for the hypothesis that would give L^2 approximation; the constants c , C of its Theorem III.34.1 are not explicit.

III.5 *Quantitative rounding and derandomization* (03 Research question 8): the group size needed by the sharpness constructions, and an efficient deterministic rounding. [6 October 2026; sources 06–10] Also 08’s sixth question: a deterministic construction with rounding error $o(\eta^4)$ in U^d , not only against single characters. [6 October 2026; sources 21–27] Source 21 (its Question 9) asks the least dimension needed to realize its exponent-five model by sets through a given order, and for deterministic rounding.

- III.6** *Counting that rectifies automatically* (03 Research question 6): a localized version of Theorem III.1.2 that counts ordinary integer progressions with all boundary errors. The lifting step itself is formal (Lemma III.4.5); its interval hypothesis is what a modular norm does not give. [6 October 2026; sources 15–20] Source 20 (its Question 9) asks for a localization keeping two small balanced norms in every error term.
- III.7** *Progression counting with more of the mixed information* (04 Q10): mixed norms, further cancellation, pairwise information, and non-admissible slopes with explicit subgroup averages. [6 October 2026; sources 06–10] *Advanced by source 10*: its Theorem III.28.2 needs only one invertible coefficient difference (next to the target), pays an explicit kernel factor $K^{1/2^{k-1}}$ for the others, and uses mixed L^p norms; the factor is sharp on $(\mathbb{Z}/2\mathbb{Z})^n$ for three terms (Proposition III.28.4). Explicit subgroup averages, and the pairwise information, remain open.
- Added in the second write, 6 October 2026 (sources 06–10):*
- III.8** *The optimal density factor in progression counts* (06 Question 6). For $k = 3$ the variance factor of Proposition III.5.6 is sharp (singletons); for $k \geq 4$, the best coefficient in place of $\sqrt{\delta} B_k(\delta)$, first within repeated telescoping and mixed-norm Cauchy–Schwarz proofs. Source 10’s mixed-norm bound $\delta S_k(\delta)\eta$ (Corollary III.29.2) is the corresponding linear statement. [6 October 2026; sources 15–20] *Advanced by source 20*: its Corollary III.45.4 has a coefficient strictly smaller than source 06’s $\sqrt{\delta} B_k(\delta)$ for every $k \geq 3$ (note there), not claimed sharp for $k \geq 4$. Source 20 (its Question 2) asks for the best three-term coefficient at fixed density, for weights and for indicator sets (source 06’s Proposition III.5.6 has coefficient 1 at singletons only).
- III.9** *Torsion constants in the generalized von Neumann bound* (06 Question 7; 10 Questions 1–2 and the question paragraph at the end of Section III.29). The best constant C in $|\Lambda_{\mathbf{a}}(f)| \leq C \|f_j\|_{U^{k-1}}$ for a fixed group, pattern and target, first for four-term patterns on elementary abelian 3-groups (10); the two-target inequality of source 06’s Theorem III.5.1 when slope differences have kernels (06); keeping the images cG of Lemma III.28.1 instead of G and exploiting their overlaps, exact constants for larger k and general torsion profiles, and the joint choice of telescoping order and L^p exponents (10). Source 10’s Theorem III.28.2 answers the one-target form with explicit, not in general optimal, constants. [6 October 2026; sources 15–20] *Answered for the norm-conversion step by source 20*: its Theorem III.44.1 gives the optimal constant of the directional-to-global comparison for every subgroup tuple, and its Proposition III.46.1 the sharp pair correction $\sqrt{m-1}$. Source 20 asks for the full multilinear operator norm (its Question 4) and for a quotient-adapted baseline removing large torsion costs (its Question 8).
- III.10** *A correct global lower bound for the cube excess* (08 Question 2; 09 Research question 5). The largest c with $Q_3(F) - \delta^8 \geq c\delta^4 \|F - \delta\|_{U^2}^4$ for all $0 \leq F \leq 1$ on odd-order groups; $2 \leq c \leq 15337/1296 < 12$ (08; the upper end is Remark W1). 09 asks for natural hypotheses under which $M_5 \geq 0$, or a sharper bound on its negative part, giving a corrected lower bound. Minimizing $12 - 2s + 6s^2 + s^4$ over the two-frequency family lowers the upper end slightly; that family is not known to be extremal.
- III.11** *Computing the survivor polynomial* (08 Question 4): $\sum_r A_{d,r}(G)z^r$ faster than by enumerating all subsets, keeping small-characteristic exceptions, with certificates a proof assistant can check. [6 October 2026; sources 11–14] Source 13’s coloopless polynomial $\sum_j c_{d,p}(j)z^j$ (its Question 8) is the same object; it suggests organizing the enumeration by flats of the dual matroid, with characteristic-dependent minors kept visible.

- III.12** *Stability of quartic-coefficient extremizers* (08 Question 5; 09 Research question 4). Structure of f when $R_d E_4(f) + T_d E_{\langle 2 \rangle}(f)$ (09: $P_d E_4 + T_d T_2$) is nearly $A_d \|f\|_{U^d}^4$: concentration on one or a few characters; on odd-order groups the extremal model is not known. A different problem from item III.4. [6 October 2026; sources 15–20] Also source 16 (its Question 12: near-maximizers of the complete cube polynomial under fixed lower-order norms) and source 18 (its Question 8: stability of B_d -sharpness, separating weights from indicator sets).
- III.13** *Cube expansions localized to intervals, progressions or Bohr sets* (08 Question 7; 09 Research question 6). Which of the three exact cancellations survive when the cube is restricted, with explicit boundary terms; 09 suggests expanding around a weighted reference measure adapted to the interval.
- III.14** *Sparse majorized weights* (08 Question 8): a transfer of the inhomogeneous expansion (Theorem III.13.1) to sparse weights with an explicit error budget, expanding around the majorant model.
- III.15** *Other linear systems* (08 Question 10; 09 Research question 9). For a fixed system of integer linear forms, the least support size at which a centered term survives and the first torsion-sensitive terms, combined with true-complexity estimates to get the correct first nonzero power; progressions need not inherit the cube's three cancellations. [6 October 2026; sources 15–20] Source 20 (its Question 11) asks for the subgroups $L_i(\ker L_j)$ of integer matrix forms through Smith normal forms or determinantal ideals.
- III.16** *The prime-onset function* (09 Research question 2): $h(7)$, $h(11)$, and the gap between (III.23.4) and $h(\ell) \leq \ell + 2$; known $h(2) = 4$, $h(3) = 5$, $h(5) = 6$. [6 October 2026; sources 34–39] Answered by source 34: $h(p) = D_{\text{div}}(p) + 1$, the least order of a nonsingular 0/1 matrix whose determinant is divisible by p (Theorem III.75.2); $h(7) = 7$, $h(11) = h(13) = h(17) = 8$ (Theorem III.76.2); $h(p) \sim h_{\text{cyc}}(p) \sim 2 \log p / \log \log p$, so the gap to $h(\ell) \leq \ell + 2$ is closed to leading order (Theorem III.79.1); universal positivity costs a factor two (Section III.80). Still asked (34 Questions 1–2): is $h(p) = h_{\text{cyc}}(p)$ for every prime, and the second-order term (coefficient between 1 and 3).
- III.17** *The constrained extremal envelope* (09 Research question 3): the largest $\mathcal{C}_d(F)$ at fixed d , δ , $\|F - \delta\|_{U^d}$, $\|F - \delta\|_{U^2}$ and $\max |F - \delta|$. [6 October 2026; sources 11–14] Source 13 (its Question 9) asks the same at $d = 3$ with $\delta, \rho, \eta, \sigma$ fixed, and which values of T_5, B_1, B_2, B_3, T_7 are jointly realizable (Theorem III.31.5). [6 October 2026; sources 15–20] Source 16 (its Question 11) asks for sharp remainder estimates in terms of lower Gowers norms, the density and the torsion profile. [6 October 2026; sources 34–39] Source 34 (its Question 10) asks the same at $d = 4$: $\sup Q_4(F)$ at fixed density and fixed $\|F - \delta\|_{U^4} = u$, bounding the quartic contribution and the remainder jointly through their Fourier phases. [7 October 2026; sources 60–69] Also asked by source 63 (its Research question 7.11): the optimal cube-count envelope when density, dimension and $\|f\|_{U^d}$ vary together.
- III.18** *A residual-sensitive inverse step* (09 Research question 7): whether the sixth-order certificate (III.21.11) and estimates of the four Fourier statistics improve a concrete inverse or counting algorithm, with a full cost analysis.
- III.19** *Polynomial colourings of higher degree* (10 Question 8). For $D_\ell > k$ only the bound of Proposition III.30.5 is known; for $k = 4$ and a scalar degree 5 or 6 the largest dimension lies between 4 and 8; start with dimension five. Proposed by source 10, not an established

conjecture. [6 October 2026; sources 11–14] Source 14 (its Question 7) restates this problem (scalar degree five, $k = 4$, $4 \leq d \leq 8$, beginning with dimension five); its Proposition III.36.7 rules out the first norm extension in characteristic five but not other quintic maps. [6 October 2026; sources 21–27] Also asked by source 22 (its Question 8). *Advanced by source 22*: the pointwise budget eight is sharp at a prescribed centre ($N_3 + N_5$, Section III.62), so no argument that works centre by centre can lower it; every extension $x_0 + N_3(u) + R(w)$ by an independent block fails (Theorem III.64.2, with the exact excess in Theorem III.64.3); and norm layers of degree divisible by the characteristic carry a quantitative excess (Theorem III.65.3). The range $4 \leq d \leq 8$ is not closed. [6 October 2026; sources 34–39] Also asked by source 38 (its Question 13: can a genuinely coupled quintic construction exceed dimension four, or does a global argument force less than eight?); its six-term count does not resolve the gap. [7 October 2026; sources 60–69] *Source 67* classifies independent norm layers with no degree cap for $2m = 4, 6, 8$ and $q \geq (2m - 1)^2$ (Theorem III.109.1). It asks (its Questions 7.8–7.9): whether $M_{2m+1,m}(p) > 0$ for every m and prime $2m < p < (2m - 1)^2$; the largest n for unrestricted scalar quintics avoiding symmetric four-term progressions (this item’s quintic question).

III.20 *Stability at the dimension threshold* (10 Question 9): must an avoiding colouring at $d = \sum_{\ell} [D_{\ell}/2]^2$ resemble the odd-degree norm construction, first in the cubic four-dimensional case? [6 October 2026; sources 11–14] Source 14 (its Question 8) asks to classify the exact extremizers: all cubic maps at the threshold in dimension four and the possible centre-dependent subspaces $S_z(Q)$ above it; the example $x + y^2$ excludes a fixed kernel. [6 October 2026; sources 21–27] Source 22 (its Question 10) asks to classify all pointwise extremizers for a given evaluation matrix; its Theorem III.63.1 (all scalar quadratic minimizers) is a first instance of source 14’s Question 8.

III.21 *Actual colours against declared coordinates* (10 Question 10): can a map with more coordinates but a small image beat the full-palette exponent of Theorem III.30.3? [6 October 2026; sources 11–14] Source 14 (its Question 10) asks the same: replace the degree-vector budget by an invariant of the actual image.

III.22 *Unrestricted symmetric-progression colourings* (10 Question 11). Deng, Tidor and Zhao ask whether $N^{o(1)}$ colours suffice to avoid nontrivial symmetrically coloured four-term progressions [37]; Shi and Dong give $O(N^{1/4})$ [33]. An open question of the literature, recorded by source 10; its theorem only limits one bounded-degree polynomial model. [7 October 2026; sources 60–69] *Source 64* transfers such colourings exactly to the circle and torus (nine copies, mirror labels) and, with Shi and Dong’s prime-indexed colourings, gets $\rho_4(\alpha) \ll \alpha^5$ and $\rho_6(\alpha) \ll \alpha^{7.847}$ (Section III.100). It asks (its Question 10.4): symmetric six-progression-free colourings with $O(N^{\gamma})$ colours for some $\gamma < 1/9$, and a four-term exponent better than $1/4$.

Added in the third write, 6 October 2026 (sources 11–14):

III.23 *Beyond the secondary coefficient* (14 Questions 1–2). On groups of exponent five: does $M_{\delta,G}(u) = \delta^8 + 6\sqrt{2}\delta^4 u^4 + C_5 \delta^2 u^6 + o(u^6)$ hold, and is $C_5 = 2^{3/4}/3$ (Theorem III.35.2 gives the exponent and a lower model; the intake’s adversarial search gave 0.56062 on $\mathbb{Z}/5\mathbb{Z}$ and 0.5626 on $(\mathbb{Z}/5\mathbb{Z})^2$ against $2^{3/4}/3 = 0.56060$, numerical evidence only)? Elsewhere, the first term after $C_* \delta^{8/3} u^{16/3}$, where order seven already changes the norm formula. [6 October 2026; sources 21–27] *Source 14’s Question 1 is answered by sources 21 and 23*: the limit exists and $C_5 = 2^{3/4}/3$, uniformly in the rank, with an absolute $O(u^8)$ remainder (source 21’s Theorem F.26.1, source 23’s Theorem III.58.1). The terms beyond are item III.27. The “elsewhere” part (source 14’s Question 2) is open, and is also asked by source 21 (its Question 6: exponent-three

groups, and mixed odd torsion by the order of the largest character) and source 23 (its Question 8: mixed torsion groups, with and without conditioning on the dominant character order).

- III.24** *Second-minimum direction counts* (14 Question 9): the next attainable $|S_z(Q)|$ in dimensions $d \geq B + 2$, and over nonprime fields in dimension $B + 1$ (Proposition III.36.6 is prime-only).

Added in the fourth write, 6 October 2026 (sources 15–20):

- III.25** *The exponent of symmetric counting* (20 Questions 1 and 3). For fixed $k \geq 4$ and density bounded away from 0 and 1, the largest θ with $|\Lambda_k(1_A) - \delta^k| \leq C_{k,\delta} u^\theta$ uniformly over admissible groups ($\theta \geq 2$ is proved; optimal for $k = 3$ on $\mathbb{F}_{29}^{2m}$, Theorem III.48.3); and a prime-cyclic analogue of the sharpness construction.

- III.26** *Equality, stability and extensions of the two-function argument* (20 Questions 5–7): equality in Theorem III.43.3 beyond constants and stability near it; an exact optimization of second-order elimination with unequal densities; systems admitting two uniformity norms and several L^p factors.

Added in the fifth write, 6 October 2026 (sources 21–27):

- III.27** *The eighth-order term in exponent five* (21 Questions 1–5 and 8; 23 Questions 2 and 7). Does the universal coefficient of u^8 exist, and is it $20/9$ (21, 23: Theorem F.26.1 bounds it, Corollary III.55.2 gives at least $20/9$; the intake’s search on $(\mathbb{Z}/5\mathbb{Z})^2$ found nothing above the lifted model, evidence only)? Is it $20/9$ on every fixed $\mathbb{F}_5^n$, with extremizers depending on one functional? A quantitative eighth-order stability theorem, and explicit bounds turning the deficit below $2^{3/4}/3$ into phase, amplitude and tail errors (21, 23). The u^{10} and u^{12} terms on $\mathbb{Z}/5\mathbb{Z}$ and how far the analytic maximizing branches continue (21). Optimal constants and simultaneous equality cases of the profile inequalities of Theorem III.51.2 (21). [6 October 2026; sources 34–39] Answered by source 38 in the stated ranges: for $u/\delta \leq 10^{-5}$ the exponent-five maximum reduces exactly to $\mathbb{Z}/5\mathbb{Z}$ on every group of exponent five, every maximizer factoring through a homomorphism onto $\mathbb{Z}/5\mathbb{Z}$, so the universal u^8 coefficient exists and is $20/9$ on every $\mathbb{F}_5^n$ (Theorem III.84.1); the u^{10} and u^{12} coefficients are $869/(216 \cdot 2^{3/4})$ and $65759/(5832 \cdot 2^{3/2})$ (Theorem III.85.1), and for every fixed odd exponent m there is a cyclic reduction for $u/\delta \leq 10^{-8}m^{-8}$ (Theorem III.86.1); quantitative off-quotient control (Corollary III.84.2) advances the stability question. Source 38’s own further questions are item III.30.

- III.28** *Pointwise and global pattern minima* (22 Questions 9 and 11–14). For which patterns and degree caps is $q^{\max(n-B,0)}$ attained by one polynomial at every centre; surjectivity of reflection maps beyond complete norm cores, and lower bounds for every fibre; defect formulas for five or more odd norm layers; stability of the zero-defect classification (Corollary III.64.4); norm lines in degrees not divisible by the characteristic. [6 October 2026; sources 34–39] Advanced by source 38 for the linear–cubic–quintic core: an exact six-term count with a quadratic-character correction and a zero-defect classification (Theorem III.88.2). Source 38 asks (its Questions 11–12) for the eight-term and higher defects (degrees 1, 3, 5, 7) and a stability theorem for small extension defect. [7 October 2026; sources 60–69] Also asked by source 67 (its Question 7.10): robustness of the norm-layer obstruction under cross-block interactions, and the interactions forcing positive supersaturation.

Added in the seventh write, 6 October 2026 (sources 34–39):

- III.29** *Binary determinants, cyclic detectors and positive supports* (34 Questions 3–5). Can every integer up to $\exp(\frac{1}{2}N \log N - o(N \log \log N))$ be the sole nonunit Smith factor of a binary matrix of order N , or at least can the coefficient of $N \log \log N$ be pushed below one (Theorem III.78.1 has one, through Shitov’s composition)? The exact $h_{\text{pos}}(m)$ over a substantial range, when $h_{\text{pos}}(m) \leq 2h_{\text{cyc}}(m)$ is strict, and its second-order term; the least affinely independent Boolean configuration with a prescribed finite abelian row quotient and all coordinate generators nonzero.
- III.30** *Quotient reduction in exponent five and odd exponent* (38 Questions 1–6). The supremum c_5^* of the ranges $u/\delta < c$ with exact five-point reduction (proved: $c_5^* \geq 10^{-5}$) and what first destroys it; whether the odd-exponent cutoff $10^{-8}m^{-8}$ must depend on m , or a sequence of noncyclic extremizers with parameters tending to zero; the cyclic profiles $\mathcal{M}_{1,\mathbb{Z}/m\mathbb{Z}}(t)$ for small odd m and their arithmetic transitions; an exact algebraic description of the five-point branch and the first nonzero term of $\alpha(t) - \pi/10$; optimal constants in the quotient-stability bounds (the factor 256 is not optimized); the constrained weighted profile once the five-point function reaches 0 or 1, and how fast indicator sets approximate it.
- III.31** *Binomial separation and endpoint exponents at Gowers’s Section 4* (64 Questions 10.1–10.3 and 10.5–10.8). [7 October 2026; sources 60–69] Is $\beta_6 = 3$, with $(1, 5, 10)$ -separating sets in intervals of length $r^{3+o(1)}$? Mirror-separating six-binomial sets modulo some $m \leq Cr^3$; the cyclic capacity $\sup \log |D|/\log B$ (is it $1/3$, never attained?); whether $\rho_4(\alpha)$ or $\rho_6(\alpha)$ decays faster than every fixed power; other even lengths and symmetric patterns, and non-mirror trivial partitions; practical bounds $P_0(\alpha, \eta, k)$ from the prime-modulus realization; fewer than nine palette copies, and a formal check of the six-variable chain. Section III.100; the colouring question is in item III.22.

The seven items carry the eight questions: 03 Research questions 1, 2, 6 and 8, and 04 Q1–Q3 and Q10. [6 October 2026; sources 06–10] With sources 06–10 the twenty-two items carry thirty-three questions and one question paragraph: also 06 Questions 6–7; 08 Questions 1–8 and 10; 09 Research questions 1–7 and 9; 10 Questions 1–2 and 8–11, and the paragraph at the end of 10’s Section III.29. [6 October 2026; sources 11–14] With sources 11–14 the twenty-four items carry forty-four questions and one question paragraph: also 13 Questions 8–10 and 14 Questions 1–3 and 7–11. [6 October 2026; sources 15–20] With sources 15–20 the twenty-six items carry sixty questions and one question paragraph: also 16 Questions 10–12, 18 Questions 6–8, and 20 Questions 1–9 and 11. [6 October 2026; sources 21–27] With sources 21–27 the twenty-eight items carry seventy-nine questions and one question paragraph: also 21 Questions 1–9, 22 Questions 8–14, and 23 Questions 2, 7 and 8. [6 October 2026; sources 28–33] With sources 28–33 the twenty-eight items carry eighty-eight questions and one question paragraph: also 28 Questions 7–11 and 29 Questions 1–4. [6 October 2026; sources 34–39] With sources 34–39 the thirty items carry one hundred and eleven questions and one question paragraph: also 34 Questions 1–10, 35 Questions 1 and 7–9, and 38 Questions 1–6 and 11–13. [7 October 2026; sources 40–46] With sources 40–46 the thirty items carry one hundred and fifteen questions and one question paragraph: also 45 Questions 6.11–6.14. [7 October 2026; sources 47–53] With sources 47–53 the thirty items carry one hundred and twenty questions and one question paragraph: also 49 Questions 10.8–10.12. [7 October 2026; sources 60–69] With sources 60–69 the thirty-one items carry one hundred and forty-one questions and one question paragraph: also 60 Research questions 1–4, 63 Research questions 7.1–7.5 and 7.11, 64 Questions 10.1–10.8 and 67 Questions 7.8–7.10.

Theme IV: floor patterns (7 questions)

- IV.1** *The second-order carry entropy* (01 Question 1; the remainder of 04 Q8). Is there a $-k \log_2 k$ term in $\log_2 C_k$, or is $\log_2 C_k = k^2 + O(k)$? Theorem IV.2.3 leaves exactly this scale; counting the hyperplane intersections that meet the open cube is the suggested route. The leading constant 1 is settled (dated note at 04's Q8). [6 October 2026; sources 06–10] Source 06's Question 2 asks the same; its Theorem IV.7.6 settles the leading constant independently and uniformly modulo M .
- IV.2** *Exact counts, arrangement structure, certified C_5 and A_4* (01 Question 2; 04 Q8, its computational part). A recurrence, an intersection-poset formula or a symmetry-reduced enumeration for C_k and A_k (01); certified counts for $k = 4$ (04): $C_4 = F_4 = 154$ is certified by source 01, A_4 is not. *Intake values, not certified:* $C_5 = 8410$ and $A_4 = 770$, from value-branching floating-point linear programs with every feasibility witness re-evaluated exactly. The infeasibility branches rest on solver status, not on rational certificates. Missing: rational certificates, for instance from source 01's or source 04's certificate generator, whose formats both handle these dimensions in principle. The values satisfy the proved bounds (note in Section IV.6). [6 October 2026; sources 06–10] Source 06's Question 1 asks for a structural formula or recurrence for C_k and D_k (its D_k is the A_k here), from the certified $C_4 = 154$ and $D_3 = 38$, and warns that an OEIS match must identify the objects, not a prefix. Source 06 certifies $C_4 = 154$ a second time; it does not compute C_5 or $D_4 = A_4$.
- IV.3** *The modulus dependence and the realizable profiles* (01 Question 3; 04 Q9). When modular collisions stabilize, and the exact number of translation classes as a function of M (01); which global and fixed-slope profiles actually occur for slopes from a short discrete interval, whether an average multiplicity can replace the worst case 2^k , and small moduli (04). [6 October 2026; sources 06–10] Also source 06's Questions 3 and 4: the ratio $C_{k,M}/C_k$ (or its logarithmic scale) for $M = 2, 3$, and an exact formula for the number $L_k(r, M)$ of affine translation classes in the saturated range $2r \geq M$.
- IV.4** *Reducing the selection loss* (01 Question 4). Which orthogonality or positivity hypotheses allow a phase-removal interface to lose L^{-1} instead of L^{-2} ; for arbitrary complex summands the L^{-2} loss is sharp. Source 04's Lemma IV.6.6 already replaces B_k^2 by $2^k B_k$ through local multiplicity. [6 October 2026; sources 06–10] Source 06's Question 5 asks the same for its Theorem IV.8.1: an orthogonality or energy-weighted decomposition that pays only for the classes carrying energy.
- IV.5** *Polynomial floor entropy beyond the leading term* (01 Question 9): secondary terms of $\log_2 |\mathcal{P}_{n,d}(r)|$ and the transition when $\log r \asymp n$.
Added in the second write, 6 October 2026 (sources 06–10):
- IV.6** *Carry entropy of other coefficient families* (06 Question 11): which finite families $V \subset \mathbb{Z}_{\geq 0}^k$ have carry entropy comparable to rank times the logarithm of their size, and, for signed families, the lower-dimensional floor-boundary obstruction first (source 06's remark after Theorem IV.7.3; source 01's example before Section IV.4).
Added in the fourth write, 6 October 2026 (sources 15–20):
- IV.7** *The best local phase-removal coefficient* (16 Questions 1–4). The largest $c_k(C)$ for cells between $m/(3k)$ and Ckm (the $k \log k$ order is settled by Theorem IV.13.2 and Proposition IV.13.5; a

factor exponential in k remains); a coefficient $2^{-O(k \log k)}$ with cells of size $O(m)$ (source 04's Theorem IV.6.7 keeps short cells at $2^{-k^2+O(k)}$, source 16's Proposition IV.14.2 at a global count); nonuniform windows keeping all weights; localization for nonclassical torsion-valued phases. Related to item IV.4. [7 October 2026; sources 60–69] Answered to first order by source 63: a positive face cover retains v_k/D_k with $D_k \leq e^{2/(3(k+2))}$, cells of lengths $\lceil m/(3k) \rceil$ or one more, so the best coefficient is asymptotic to $v_k = 2^{k+1}/(k+2)!$ (Theorem IV.68.1); the fixed-degree optimum remains. Its cells are at most m long, which bears on the second question (not checked uniformly in m here). Source 63 asks (its Research questions 7.6–7.10): the exact positive-cover cost; the next term $\lim_k k(1 - c_k^{\text{loc}}/v_k)$; c_k^{loc} at $k = 1, 2$; lattice corrections and simultaneous limits; other windows, boxes and kernels.

Added in the seventh write, 6 October 2026 (sources 34–39):

IV.8 *Symmetry certification beyond the threshold theorem* (36 Questions 1–6, 10 and 11). All endpoint configurations (f, σ) with $\mathcal{A}_f(\sigma) = \vartheta_{k,p}$ beyond products of linear functionals; stability at the threshold, in a metric to be chosen first; a joint rank invariant for simultaneous alternating defects; below-threshold coherent rank reduction for one-function derivative arrays; versions localized to progressions or Bohr sets with explicit boundary errors that would refine Proposition 17.7 (Section IV.10 has source 16's local extraction); removal of lower-order multiaffine components; certification from a coefficient tensor and sampled derivative coefficients with a robust margin; inverse arguments in which the selector arrives with energy a and $p > k/a$ without destroying the advantage.

IV.9 *Small characteristic and composite moduli* (36 Questions 7–9). Repeated-variable identities forced near energy one in small characteristic, so that a sharp threshold certifies the hypotheses of a nonclassical integration theorem; optimal bilinear factors from the elementary divisors of the alternating bicharacter over odd composite exponent, and an exact threshold over $\mathbb{Z}/m\mathbb{Z}$; an explicit characteristic-two version, deciding which parts keep the threshold $1 - 2^{-k}$. Sources 43 and 46, placed and not yet written, record characteristic-two results (for their writes). [7 October 2026; sources 40–46] Answered for symmetric trilinear forms over $\mathbb{F}_2^n$ by sources 46 and 43 (exact energy $2^{-r/2}$, codimension $r/2$; Section IV.25); over $\mathbb{F}_2$ the threshold $1 - 2^{1-d}$ certifies nonclassical integrability in every degree (source 43's Theorem IV.38.4); on every finite abelian group the orthogonality factor d_B^{-1} and the symmetry threshold $1 - (1 - 1/p_*(G))^k$ (Theorems IV.35.2, IV.36.4). Still asked: a spectral identity detecting the p -repetition conditions in odd small characteristic (43 Q5); integration on composite groups from elementary divisors (43 Q6); a sharp mixed cube inequality controlled by rank C_T over $\mathbb{F}_p$ at degree $p+1$ (46 Q6). [7 October 2026; sources 47–53] Source 47 asks (its Research question 7) for the analogue of its forbidden-frequency argument in characteristic $p > 2$, keeping the p -repetition transfer conditions of nonclassical integration and the full translation action on repeated-direction products. [7 October 2026; sources 60–69] Also asked by sources 65 and 69 (65 Research question 6, 69 Research question 7): a contraction radical for the p -repetition defect in odd characteristic, and the analogue of the threshold.

Added in the eighth write, 7 October 2026 (sources 40–46):

IV.10 *Binary cubic norms and Boolean integration* (43 Questions 1–4 and 7–9; 46 Questions 1–5 and 7–9). Whether the quartic extremum stabilizes, $\mathcal{M}(C_{4,n}) = 11/16$ for every $n \geq 2$ (first $n = 3$; proved: $11/16 \leq M_{4,n} \leq 3/4$); the exact seven-function supremum for a symmetric binary trilinear form, whether it depends only on rank C_T , and whether the exponent $r/4$ and the codimension coefficient two are optimal (43 Q2, 46 Q1); all extremizers of the cubic norm,

and stability from near-extremal energy to an exact extremizer in a dimension-free way (43 Q3, 46 Q2–Q3); the families $z \mapsto \Omega_z$ that arise from symmetric d -linear tensors for $d \geq 5$; the ordinary slice-rank repair number $\rho(T)$ between $r/2$ and r ; the extremum over a prescribed root-of-unity alphabet (fourth roots give $7/16 < 1/2$ on $\mathbb{F}_4$); the largest integrable subspaces of the trace tensors for $d \geq p + 2$, and Frobenius-twisted trace families; robustness before exact symmetry, with each loss tracked; a boundary-sensitive local twirl on boxes or Bohr sets; an inverse-theorem regime producing a selector above the certification threshold; removal of multiaffine offsets, given the constant and lower-order terms of `IsMultilinear`. Section IV.25; the odd-characteristic and composite questions are merged into item IV.9. Source 47, placed by `f8bc5e2ec` and not yet written in, records the maximum $1 - (d + 1)/2^d$ for C_d in every degree and dimension, which would answer the first question (for its write). [7 October 2026; sources 47–53] *The first question is answered by source 47: $\mathcal{M}(C_{4,n}) = 11/16$ for every $n \geq 2$, the canonical maximum being $1 - (d + 1)/2^d$ in every degree, and $11/16$ is the exact symmetric quartic threshold (Theorems F.52.1, F.52.2); source 47’s questions are item IV.11.*

Added in the ninth write, 7 October 2026 (sources 47–53):

- IV.11** *Boolean obstruction energies beyond degree six* (47 Research questions 1–6 and 8). The septic threshold, $120/128 \leq M_{7,n}^{\text{sym}} \leq 121/128$ (is it $15/16$?); whether $1 - (d + 1)/2^d$ is the universal symmetric extremum in every degree; the least support of a non-pure symmetric s -linear map for $s \geq 4$ ($u^{\otimes s} + v^{\otimes s}$ gives $2^{1-s} - 2^{1-2s}$ against the bound $7/2^{s+2}$); the tensors and functions attaining the quintic and sextic thresholds, and dimension-free phase stability from the deficit recursion; pullbacks along $V \twoheadrightarrow V_0$ that preserve the maximum (settled for C_d); direct sums of canonical tensors, and mixed functions on the cube vertices; quantified period and averaging errors on Bohr sets or progressions, propagated through a genuine selector construction. Section IV.42; the odd-characteristic question is in item IV.9, the formalization question in item V.2. [7 October 2026; sources 60–69] *Research questions 1, 2 and the tensor part of 4 are answered by sources 65 and 69: the maximum is $1 - (d + 1)/2^d$ in every degree (the septic value is $15/16$), with canonical endpoint tensors (Section IV.51). Asked: the second extremal value and whether $1 - 3d/2^{d+1}$ is attained (65 Research question 1, 69 Research question 1); the extremum at prescribed radical codimension (65 Research question 2); the maximizing functions and phase stability (65 Research question 3, 69 Research questions 2–3; the function part of 47’s Question 4); inactive coordinates outside the canonical class (65 Research question 4, 69 Research question 4); direct sums (65 Research question 5); local measures and global bookkeeping (65 Research question 7, 69 Research question 9); multiaffine, mixed and incompletely symmetric selectors (65 Research question 8, 69 Research question 8); energy against model size and the maximal model tensors (69 Research questions 5–6).*

The five items carry the seven questions: 01 Questions 1–4 and 9, and 04 Q8–Q9. [6 October 2026; sources 06–10] With sources 06–10 the six items carry thirteen questions: also 06 Questions 1–5 and 11. [6 October 2026; sources 15–20] With sources 15–20 the seven items carry seventeen questions: also 16 Questions 1–4. [6 October 2026; sources 34–39] With sources 34–39 the nine items carry twenty-eight questions: also 36 Questions 1–11. [7 October 2026; sources 40–46] With sources 40–46 the ten items carry forty-six questions: also 43 Questions 1–9 and 46 Questions 1–9. [7 October 2026; sources 47–53] With sources 47–53 the eleven items carry fifty-four questions: also 47 Research questions 1–8. [7 October 2026; sources 60–69] With sources 60–69 the eleven items carry seventy-six questions: also 63 Research questions 7.6–7.10, 65 Research questions 1–8 and 69 Research questions 1–9.

Theme V: integration and formalization (8 questions)

V.1 *A complete, audited propagation through one argument* (01 Question 10; 02 Subsection 11.4; 03 Research question 7; 04 Q11). All four sources ask for a dependency graph of the quantitative parameters, into which the local improvements are substituted. Every later threshold would then be recomputed under the same hypotheses, ending in an explicit final bound, compared with the original argument and with Leng–Sah–Sawhney. Source 02 proposes the length-four proof as a first target; source 03’s Proposition V.12.1 supplies the final bookkeeping. Source 01 adds that the issue flagged at Corollary 16.11 must be resolved first. That issue is still open at HEAD (`corollary_16_11`, `theorem_18_1`). [6 October 2026; sources 06–10] Also asked by source 06 (Question 10: close the structural estimates around Sections 13 and 16 and Theorem 18.1 first, keeping the local gains as parameters), source 07 (Question 8: which bottleneck dominates after its exactification, in a fully tracked bound) and source 09 (Research question 10: locate a step whose final bound actually depends on the cube excess). At HEAD of the second write `corollary_16_11`, `theorem_18_1` and `theorem_18_2` are still open. [6 October 2026; sources 11–14] Also asked by source 12 (its Questions 10–11: reprove the later local-model inputs of Sections 13 and 16 with its Theorem II.40.2, and a verified parameter calculus), source 13 (its Question 11: the best complete bound of the existing architecture after the linear and real-valued replacements) and source 14 (its Question 12: insert its profile and geometric denominators into a whole extraction chain). [6 October 2026; sources 15–20] Also asked by source 15 (its Question 4: how much of the coupled gain survives Sections 15–18, through a parameter ledger carrying $(\beta, \gamma, \chi, \rho_3)$), source 16 (its Question 13: substitute the new cell scale and recurrence thresholds), source 18 (its Question 12: a checked Lemma 15.5 threshold and a checked conditional partition transfer first), source 19 (its Question 12) and source 20 (its Question 10). At HEAD 0b4890a1c `lemma_16_10`, `corollary_16_11`, `theorem_18_1` and `theorem_18_2` are still open. [6 October 2026; sources 21–27] Also asked by source 21 (its Question 11: identify a downstream theorem that uses the characteristic-five cube count before claiming an improvement), source 22 (its Question 7), source 25 (its Question 8) and source 27 (its Question 11: a complete Section 15 threshold audit through Lemma 15.6). At HEAD 2dfe46f44 `lemma_16_10`, `corollary_16_11`, `theorem_18_1` and `theorem_18_2` are still open. [6 October 2026; sources 28–33] Also asked by source 28 (its Questions 12, 14 and 15: propagate through Section 13, an end-to-end explicit density increment, and a bridge to modern inverse theorems), source 29 (its Question 13: a complete quantitative substitution), source 30 (its Questions 8–9: propagate its radius through Theorem 10.13 and find which losses dominate), source 31 (its Question 13: a checked quantitative application at one progression length) and source 33 (its Question 12: actual propagation through the Gowers pipeline). At HEAD 5d4718772 `lemma_16_10`, `corollary_16_11`, `theorem_18_1` and `theorem_18_2` are still open. [6 October 2026; sources 34–39] Also asked by source 34 (its Question 11: recompute a normalized scheme with its cube bound and rigidity threshold), source 35 (its Question 13: the graded selection inserted into the whole corrected argument), source 37 (its Question 7: rework the degree induction with the short-step estimate, then propagate only verified improvements) and source 38 (its Question 15: a density-increment or transference argument in which one of its local statements survives). `lemma_16_10`, `corollary_16_11`, `theorem_18_1` and `theorem_18_2` are still open at HEAD 89c1cbc38. [7 October 2026; sources 40–46] Also asked by source 40 (its Question 6: the bounds that result from substituting its Corollary II.170.1 into the later steps, with a parameter ledger), source 44 (its Question 9: the complete Section 15 propagation of its principal term and incidence coefficient) and source 45 (its Question 6.16: a complete argument in which one

of its local gains survives). `corollary_5_8` is now exact; `corollary_16_11`, `theorem_18_1` and `theorem_18_2` are still open at HEAD 6d4942222. [7 October 2026; sources 47–53] Also asked by source 48 (its Research question 10: which later uses of the restriction can consume the constructed subset without infeasible enumeration, tracking density losses and execution costs), source 49 (its Question 10.15: carry its Section 16 controls through the later phase-extraction and density-increment arguments, assuming I1–I7 certified) and source 53 (its item 4: compose its actual-density refinement with the Fourier and global iteration interfaces, preserving their hypotheses and printed exponents). `lemma_16_10` (refuted in its catalogue form), `corollary_16_11`, `theorem_18_1` and `theorem_18_2` are still open at HEAD 06b30ca52. [7 October 2026; sources 54–59] Also asked by source 54 (its Question 11.10: propagate the common-phase theorem through Section 8 with every normalization and loss), source 56 (its Question 12.7: relative repair at every high-agreement stage of a full inverse argument), source 57 (its Research questions 6 and 11: Fejér selection through the Section 16 induction toward Corollary 16.11, and a complete density-increment iteration from the revised local estimates) and source 58 (its Question 13.9: the new restriction constants through the later stages). `theorem_13_12` is exact at HEAD fe7165a3c; `lemma_16_10`, `corollary_16_11`, `theorem_18_1` and `theorem_18_2` are still open. [7 October 2026; sources 60–69] Also asked: a complete propagation of the local improvements (60 Research question 14; 62 Question 12.8; 63 Research question 7.12; 67 Question 7.11).

V.2 *Formal libraries, and formalizing sharpness* (02 Subsection 11.7; 03 Research question 9; 04 Q12; 05 paragraph 7). A characteristic-independent library: weighted Fourier energy on $G \times \widehat{G}$ and real phase freezing (02); integer linear forms, centered multilinear averages and character constraints with torsion data (03); the low-risk interfaces first, namely moment interpolation, the telescope and the support enumeration, then the selection theorem (04); and sharpness as well as existence, starting with the parity obstruction for $1/3$ and the residual inequality (I.19.4) (05). The write adds one concrete gap. Source 02’s Theorem I.11.1 could become formal now that `corollary_5_6_holds` and `polynomialOn_affine_pullback` exist; what is missing is the transfer of cells back to the input progressions and the selection step. The lifting lemma (Lemma III.4.5) is already formal. [6 October 2026; sources 06–10] Also source 08 (Question 9: the elementary fourth-order bound and the quartic torsion formula as separate verified modules, with the unnormalized target (V.32.1)) and source 09 (Research question 8: a library of reindexing equivalences and integer-lattice certificates for Boolean affine forms, keeping non-prime groups). [6 October 2026; sources 11–14] Also source 11 (its Question 9: the support-pattern and colour-comparison layer as a Mathlib development connected to the Section 6 statements) and source 13 (its Question 12: the finite sharp statements and their extremizers as reusable Lean lemmas, beginning with weighted selection). [6 October 2026; sources 15–20] Also source 15 (its Question 9: the division-free interpolation and the collision-parametrized product step), source 16 (its Question 14: the certificate theorem and the window geometry first), source 17 (its Question 10: choose L by a verified finite search), source 19 (its Question 11: the hypergeometric identity, the marked-sample bound, the minimax construction and the interpolation theorem) and source 20 (its Question 12: the two-function theorem, the index comparison and the second-order telescope, root-free). [6 October 2026; sources 21–27] Also source 21 (its Question 10: the 34-monomial five-point certificate and the sextic support identities by a small trusted evaluator), source 22 (its Question 15: $J_{k,2}$, $\delta_{1,1}(p)$, the weighted cumulative-sum lemma and the evaluation-matrix minimum), source 23 (its Question 12: the semigroup and mixed-radix lower bounds first), source 24 (its Question 10: the selector, walk injection, fibre compression and compiler over general finite sets), source 26 (its Question 5:

the pair-counting, residue-partition and signed-mass lemmas, composed into the $k = 1$ instance of `corollary_5_8`) and source 27 (its Question 12: the first-order theorem, then $d = 2$, then the exact count with a verified rational-incidence checker). [6 October 2026; sources 28–33] Also source 30 (its Question 12: the escape theorem, the packing theorem and the margin-preserving extension, with endpoint-safe non-return and opposite-sign tube disjointness as the main obligations) and source 33 (its Question 11: the auxiliary-label correction, affine separation, proper-coordinate independence and the one-error identity, starting from the rational inequality $\varepsilon_s \geq s\delta - 2s(s-1)\delta^2$). [6 October 2026; sources 34–39] Also source 35 (its Question 15: one library for weighted configurations, so that later refinements change only a rank estimate, a coefficient mass or a finite energy identity), source 36 (its Question 12: the eigenspace-multiplicity proof from finite-dimensional spectral theory, connected to the global phase-extraction predicate), source 37 (its Question 9: certified algebraic extremizers, first the lower-bound track), source 38 (its Question 14: the centered cube coefficients, energy coercivity, the field-fibre count and the Weyl-coset expansion) and source 39 (its Question 7: the sampler at `Section16SampledOrAnchoredOn`, then labelled hereditary refinement with a dependent list type). [7 October 2026; sources 40–46] Also source 40 (its Question 8: the one-feature relation theorem, the positive-kernel moment and the score selection as reusable components, with the finite checks generated from the same definitions), source 41 (its item 5: the finite-group mass and variance lemmas, residue partitions, explicit recurrence and numerical induction composed into the pinned predicate; the repository has since formalized the whole interface by its own route, Section V.1), source 43 (its Question 10: the projective eigenvalue-multiplicity argument, the Boolean primitive formula and a small proof object for the quartic value), source 44 (its Question 8: the smallest trusted checker for the eight-shape formula and $b_{2,8}$), source 45 (its Question 6.15: the energy identity, the unit-Schur-defect lemma, the complement identity and the rational U^4 certificate) and source 46 (its Question 10: certificates for the symplectic decomposition, the residual integrator and the extremal value). [7 October 2026; sources 47–53] Also source 47 (its Research question 9: the even-period and fixed-derivative lemmas, then the canonical induction, with a finite-linear-algebra module for the pure-defect classification and the support census), source 48 (its Research question 9: the insertion and response invariants and a compact certificate with a cheap checker), source 49 (its Question 10.16: the translation reduction and rank-two classification, the two finite certificates, and the cover recurrence as a dimension-indexed expression type), source 50 (its Research question 10: the negative theorem, and a catalogue separating the printed lemma, the false abstraction and a replacement; the repository has since formalized the negation) and source 52 (the negations of the two preserved Stage 13.7 encodings; not formalized at HEAD). [7 October 2026; sources 54–59] Also source 54 (its Question 11.11: a finite-abelian-group quadratic-refinement library with doubled roots), source 55 (its Question 12.10: the four-vertex classification, the six-triple Hölder estimate and the explicit local theorem, with the entry theorem as a named interface), source 56 (its Question 12.10: an abstract domain certificate, formalized once), source 57 (its Research question 12: the purification-to-square synthesis and the near-maximal endpoint, discharging `theorem_13_12`, done in the repository’s own form since, and the reflection-row lemma and partition classification, not done) and source 58 (its Question 13.10: the expectation theorem once for all finite abelian groups, with numerical corollaries by exact certificates). [7 October 2026; sources 60–69] Also asked: formalizing the bounded inequality, the face-cover identity and the finite parameter theorem (63 Research question 7.13); certified reconstruction and endpoint recognition (65 Research question 9); verified finite certificates for the minimal quotient (69 Research question 10); the partition normal form and the four levels with a certificate generator (66 Question 13.10).

Added in the second write, 6 October 2026 (sources 06–10):

V.3 *Finite corrections to the iteration potential* (06 Question 9). The leading cost of Theorem V.25.2 is optimal for its abstract rule; asked are a minimal discrete correction or an exact dynamic-programming potential with a usable closed form, and the treatment of density-dependent length prefactors and nonmonotone thresholds without worst-case replacement.

The two items carry the eight questions: 01 Question 10, 02 Subsections 11.4 and 11.7, 03 Research questions 7 and 9, 04 Q11–Q12, and 05 paragraph 7. [6 October 2026; sources 06–10] With sources 06–10 the three items carry fourteen questions: also 06 Questions 9–10, 07 Question 8, 08 Question 9 and 09 Research questions 8 and 10. [6 October 2026; sources 11–14] With sources 11–14 the three items carry twenty questions: also 11 Question 9, 12 Questions 10–11, 13 Questions 11–12 and 14 Question 12. [6 October 2026; sources 15–20] With sources 15–20 the three items carry thirty questions: also 15 Questions 4 and 9, 16 Questions 13–14, 17 Question 10, 18 Question 12, 19 Questions 11–12, and 20 Questions 10 and 12. [6 October 2026; sources 21–27] With sources 21–27 the three items carry forty questions: also 21 Questions 10–11, 22 Questions 7 and 15, 23 Question 12, 24 Question 10, 25 Question 8, 26 Question 5 and 27 Questions 11–12. [6 October 2026; sources 28–33] With sources 28–33 the three items carry fifty questions: also 28 Questions 12, 14 and 15, 29 Question 13, 30 Questions 8, 9 and 12, 31 Question 13 and 33 Questions 11–12. [6 October 2026; sources 34–39] With sources 34–39 the three items carry fifty-nine questions: also 34 Question 11, 35 Questions 13 and 15, 36 Question 12, 37 Questions 7 and 9, 38 Questions 14–15, and 39 Question 7. [7 October 2026; sources 40–46] With sources 40–46 the three items carry sixty-eight questions: also 40 Questions 6 and 8, 41 item 5, 43 Question 10, 44 Questions 8–9, 45 Questions 6.15–6.16, and 46 Question 10. [7 October 2026; sources 47–53] With sources 47–53 the three items carry seventy-six questions: also 47 Research question 9, 48 Research questions 9–10, 49 Questions 10.15–10.16, 50 Research question 10, one paragraph of 52 and 53 item 4. [7 October 2026; sources 54–59] With sources 54–59 the three items carry eighty-six questions: also 54 Questions 11.10–11.11, 55 Question 12.10, 56 Questions 12.7 and 12.10, 57 Research questions 6, 11 and 12, and 58 Questions 13.9–13.10. [7 October 2026; sources 60–69] With sources 60–69 the three items carry ninety-four questions: also 60 Research question 14, 62 Question 12.8, 63 Research questions 7.12–7.13, 65 Research question 9, 66 Question 13.10, 67 Question 7.11 and 69 Research question 10.

Further limitations stated by the sources (kept in place; listed here for the standing rule). Sharpness of the density transfer holds in the weighted model, not for arithmetic characters (01, Remark I.2.2). The constants 36π and $3/4$ are not claimed optimal (05). Sharpness of the simultaneous exponent is for linear phases, uniform arcs and symmetric widths (05). The coefficient 6 of the $u^{16/3}$ term is not claimed optimal, and the higher-dimensional constant $P_d/\sqrt{2}$ is not claimed sharp (04). The restriction theorem is prime-only and does not prove the Lean Lemma 9.3 or Corollary 9.4 (04). $\text{PP}(K, R_0)$ is an imported input (02). The elementary Balog–Szemerédi–Gowers bound is far weaker than Reiher–Schoen (03). R_d is not claimed to be the optimal odd-order constant (03; it is not, for $d = 3$). Finite and numerical checks prove nothing general (all five). No source has a new $r_k(N)$ bound, a priority claim, or Lean verification. [6 October 2026; sources 06–10] *Further limitations of sources 06–10* (kept in place). Source 06: the leading carry entropy rests on Zuev’s count; the ordering remark optimizes bounds, not discrepancy constants; the local phase loss is not substituted into a final theorem. Source 07: five and nine are guaranteed orders, not maximal ones; the sharpness is for the directional defect δ , not for η ; the threshold is not a replacement for an average homomorphism test; the factors 243 and 2548.03968 are not improvements of a global bound. Source 08: R_d is not claimed optimal as a multiplier of t^4 in odd order; 57 is a

convenient constant; the sets in its sharpness theorem have densities tending to δ . Source 09: no originality claim; its sixth-order certificates are not inverse theorems. Source 10: no novelty claim for the polynomial method (Führer–Taranchuk credited), the colouring theorem concerns declared polynomial coordinates only, and the Reiher–Schoen variant is labelled as using a published input. All five: finite checks prove nothing general, no Lean compilation, no new bound for $r_k(N)$. [6 October 2026; sources 11–14] *Further limitations of sources 11–14* (kept in place). Source 11: the radii $1/7$, $1/3$ and the cutoff 34 are sufficient, not optimal; “sharp” does not mean every t is a coset density; its Proposition 6.1 corollary needs $\alpha^4 > 8/9$ and does not improve the exponent four; its partial-domain theorem is not a sparse-extension claim; majority decoding and the derivative identities are prior art. Source 12: sharpness is universal, not for prime cyclic groups with the same target or for Bohr sets alone; $3\varepsilon + r\kappa$ and the factor 3 in the repair budget are not claimed optimal; the component lemmas are borrowed and not removed by the repair theorem; no claim of feasibility at $\alpha = 1/8$; Y is not claimed to contain a progression. Source 13: only the flexible-length frontier is claimed optimal; the linear specialization only; the bounded-height budget is not globally optimal; dividing by M is not optimal; the real-valued hypothesis is essential; quartic sharpness is for weights, not indicators at a fixed group size; no rectification to integers; its characteristic-two table differs. Source 14: the truncated expansion is not a finite-parameter bound; no L^2 rigidity; $2^{3/4}/3$ is a lower bound only; envelope sharpness is for an infinite family, not every ε ; its spectral bound gives no coset closeness; its symmetric colouring condition is not Corollary 8.3; the gap is prime-only; the Artin–Schreier example does not improve the bound 8; no peer review. All four: finite checks prove nothing general, no Lean compilation, no new bound for $r_k(N)$, no priority claim. *Superseded*: the first write’s “The coefficient 6 of the $u^{16/3}$ term is not claimed optimal” (source 04) stands as source 04’s statement; the optimal coefficient is $(9/4)^{1/3}$ by source 14 (Remark W4). [6 October 2026; sources 15–20] *Further limitations of sources 15–20* (kept in place). Source 15: no propagation through Sections 15–18 and no improvement of a final bound; the constant 18 and the exponents of its coupled theorem are not claimed optimal; its joint normal form classifies joint equality, not product-property extremizers; its dichotomy gives no label repair, and the saturation remark no quantitative rounding; thirteen floating-point Fourier checks are flagged as such; its questions are not claimed to be recognized open problems. Source 16: the recurrence theorem is external and its threshold not numerical; an exponential-in- k gap remains at Proposition 17.7; the own-step obstruction is for full partitions; the ceiling concerns local-scale conclusions only; the cube identities alone do not improve the inverse theorem. Source 17: prime moduli only, not a drop-in proof of the all- N lemma_9_3; thresholds not optimal or practical; the limitation concerns the scalar signed certificate only; the moment inequality is not claimed new; no equality classification for fractional weights; no efficient algorithm. Source 18: its Lemma 15.5 corollary changes one threshold, not the exponent E_k , and nothing is propagated; the results need the full degeneracy definition; Lau’s preprint is imported; the sufficient q^{-2} and necessary q^{-1} rates do not match; the three-cube circuits are classical. Source 19: C_d concerns this recovery certificate only; synchronization needs the inherited-coordinate hypothesis; Lemma 16.10 and Theorems 18.1–18.2 are not completed; Maynard’s theorem is imported and Lau’s input optional; no matching nonlinear theorem. Source 20: not a substitute for the asymmetric Corollary 3.3 under its hypotheses; B_k not claimed sharp for $k \geq 4$; indicator sharpness only on $\mathbb{F}_{29}^{2m}$; the constant c_* is crude. All six: finite checks prove nothing general, no Lean compilation, no new bound for $r_k(N)$, no priority claim. [6 October 2026; sources 21–27] *Further limitations of sources 21–27* (kept in place). Source 21: the universal u^8 coefficient is undetermined; $20/9$ is proved for the fixed group $\mathbb{Z}/5\mathbb{Z}$; indicator sharpness is a lower example; the profile constants are not optimized; the projection proposition is methodological; its review is “AI-assisted”, not peer review. Source 22: certificate class counts are not success loci; the second coefficient is not determined; the equal-allocation threshold is not

claimed to dominate; the $d = 1$ remainder is a union bound with no containment claim; global quintic avoidance is unresolved and 7,810,000 is a lower bound; “zero defect” means $\mathbb{F}_p$ -quadratic, not $\mathbb{F}_q$ -quadratic. Source 23: the Section 16 control-function comparison is not proved; the gluing is conditional on slice covers; no integer-representative fidelity; the threshold is eventual and universal, not a statement about every R ; the localization constant is sharp only for its localization problem; the tertiary cube coefficient is open. Source 24: not a new Balog–Szemerédi–Gowers record; the scalar sharpness concerns the relaxation, not realizable codegree matrices; its hashing theorem is outside its suite; its modern-input corollary imports Reiher–Schoen; nothing is propagated through Sections 8–18. Source 25: no full-arrangement census and no selection simulation; $d \geq 2$ and characteristic above $2d$ are substantive; no global improvement. Source 26: no higher-degree result; Appendix A is not the Section 16 interface; constants not optimal; the cube-root scale and the mechanisms have precedents; its builder is Linux-only. Source 27: the exact formula rests on a finite certificate not checked by a proof assistant; the cubic bound is a union bound; nothing is propagated through Lemma 15.5; polynomiality beyond eight vertices is a question; no claim of geometric irreducibility. All seven: finite checks prove nothing general, no Lean compilation, no new bound for $r_k(N)$, no priority claim. *[6 October 2026; sources 28–33] Further limitations of sources 28–33* (kept in place). Source 28: the logarithmic losses are not shown necessary; the $O(\alpha^{-3})$ count belongs to the new threshold (its old-spectrum corollary keeps the retained fraction and radius, not the count); Y depends on ϕ ; the input defect must be at most 2^{-21} ; c_4^{odd} is not determined; nothing is propagated through Sections 13, 16 or 18. Source 29: the moment inequality is sharp, the norm comparison is not asserted sharp for $d \geq 3$; the joint theorem is an absolute endpoint with row sums at most one, not unrestricted relative-energy stability; the restriction is for fixed fields and does not replace the all-moduli lemma. Source 30: the exponential base is unresolved ($\sqrt{3}$ against 3); the obstruction uses composite cyclic groups, the prime case is open; no improved Theorem 10.13 is asserted. Source 31: the qualitative critical-norm classification is Eisner and Tao’s; the joint graph inequality, the cutoff $1/120$, the odd-order envelope and the extraction exponents are not claimed optimal; Theorem II.117.1 is prime-cyclic with a fibre cap and says nothing about compatibility across slices; Appendix A is a corollary of Reiher–Schoen. Source 32: no best recurrence exponent; no parent containment in the variance branch; no higher degree; no integer rectification; the length bound exceeds one only for $N/M > 17^{39}$; its builder is Linux-only. Source 33: full or near-full domains only; F_s for $s \geq 6$ not asserted sharp; the activation thresholds are not optimal; no new bound for $r_k(N)$ or van der Waerden numbers. All six: finite checks prove nothing general, no Lean compilation, no new bound for $r_k(N)$, no priority claim. *[6 October 2026; sources 34–39] Further limitations of sources 34–39* (kept in place). Source 34: c_4^{odd} is not determined; neither endpoint is claimed optimal; the determinant range is eventual (no explicit N_0) and rests on the Alon–Vu matrices; the second-order support term is open; the transport threshold is for source 07’s finite-fibre model and is not asserted at the prime-order interface; nothing is propagated through a full iteration. Source 35: $E_{k,m}$ is a sufficient exponent, not shown optimal, and the selection’s prime-field transfer needs $p > 2(L - 1)$; the coefficient 5 is not optimal (superseded by source 34); the polynomial obstruction is necessary only, with no matching sufficient theorem in degree two or more; the prime-modulus regime $N \asymp R$ is open. Source 36: odd p only, and integration needs $p > k + 1$; for fixed p and small energy the criterion says nothing; no improvement of the scalar coefficient one; not Proposition 17.7 and not a general inverse theorem; standard Heisenberg and polarization facts are not claimed new. Source 37: the open-gap extremizers can charge endpoints and need not respect the half-open exclusion; the recurrence needs $T \geq 2^{2^{40k^3}}$ and imports the Weyl inequality; the general moment criterion is classical; no record Diophantine bound; no propagation through Corollary 5.6. Source 38: the cutoffs 10^{-5} and $10^{-8}m^{-8}$ are sufficient, not optimal; the

analytic series holds on a possibly smaller neighbourhood, not certified numerically; the envelope constant 104 is not optimal; the colouring construction is Shi and Dong's and no better colouring exponent is claimed. Source 39: minimax optimality holds for its certification rule only, and the extremizing families need not satisfy the other hypotheses of the inverse argument; occupancy methods are credited; Lemma 16.10's closing comparison is not resolved. All six: finite checks prove nothing general, no Lean compilation, no new bound for $r_k(N)$, no priority claim. [7 October 2026; sources 40–46] *Further limitations of sources 40–46* (kept in place). Source 40: the barrier is for homogeneous positive-Fourier kernels only; its field-size condition is sufficient, not optimal (source 35 improves its exponent); the small-field test is vacuous, as it says; no density-increment theorem. Source 41: the constants are deliberately weak, and the main bound is nontrivial only for $N/M > 8^{D_k}$; the variance branch may leave the cover; no rectification or optimality; it does not change the formal status. Source 42: the optimality of $1 - 1/C$ is uniform over densities (sparse extremizers), the fixed-density problem is open; no parent containment; the classical inputs are credited. Source 43: the interval $11/16 \leq M_{4,n} \leq 3/4$ is proved, constancy at $11/16$ is a conjecture; odd small characteristic and composite integration are open; the integration criterion and Heisenberg mechanism are established background. Source 44: fixed- k asymptotics, not uniform in k ; core avoidance is not full nondegeneracy; $b_{2,8}$ rests on a finite certificate. Source 45: the third energy level is not classified; the range $1/9$ is not claimed optimal; c_4^{odd} is not determined; the first energy gap is credited to inverse Pollard and Hegyvári. Source 46: the seven-function bound is not claimed sharp; the trace sections were read, not re-derived, at intake; Milićević's 2026 preprint was not audited, as it says. All seven: finite checks prove nothing general, no Lean compilation, no new bound for $r_k(N)$, no priority claim. [7 October 2026; sources 47–53] *Further limitations of sources 47–53* (kept in place). Source 47: exactness of the universal threshold only through degree six; full symmetry, the evaluation slot included, is required; $\mathbb{F}_2^n$ only, no odd-characteristic or mixed-function analogue, no pullback theorem in general. Source 48: complexity polynomial in the ambient size, not in $\log D$; dense storage; uniform fixed size is not a conditioned Bernoulli filter; no new retention exponent (the constant is source 17's). Source 49: the constants are not claimed sharp; the theorem for $k = 3$ is computer-assisted; all- k exactness is not claimed; the cover theorem is conditional on I1–I7, which it does not certify. Source 50: refutes the two-input abstraction, not Gowers's contextual lemma or Szemerédi's theorem; no Lean proof (the repository's is its own); the constants are astronomically large by design. Source 51: refutes a universal retained-data implication, not existential extraction, particular proof-term witnesses or `corollary_13_10`; its repair is existential and not compared with an independently chosen Stage 13.9 witness. Source 52: composite moduli only; the simultaneous unit-step full package is open; `lemma_13_7_holds` is untouched. Source 53: the threshold is sufficient, not optimal; no Lean implementation and no catalogue repair; the side exponent three is a convenient consequence, not optimal. All seven: finite checks prove nothing general, no new bound for $r_k(N)$, no priority claim. [7 October 2026; sources 54–59] *Further limitations of sources 54–59* (kept in place). Source 54: no lower bound on each translate; $27/4$ is not optimal for balanced indicators; repaired functions keep only the L^2 norm, not a pointwise cap; the selector of Section 8 is assumed, not constructed. Source 55: the global entry threshold is not explicit (the Eisner–Tao theorem is imported); only the first two coefficients are sharp, the radius and cubic constant are not optimized, and the exact curve is known only at $d = 2$; triple independence is essential, so Bohr sets and intervals are not covered. Source 56: the density exponent $5/4$, the thresholds $1/1024$ and the coefficient $13/50$ are not optimal; repair on a coset does not extend to the ambient group $(\mathbb{Z}/4\mathbb{Z} \rightarrow \mathbb{Z}/2\mathbb{Z})$; affine, not quadratic, repair on quadratic domains; not a stronger Lemma 10.9. Source 57: the square extraction is the repository's, Gowers's inputs are imported with their constants, the threshold N_{sq} is not made explicit in it, and the Fejér selection is source 40's; the Schur threshold $5s - 1$ is sufficient only; energy results are

torsion-free. Source 58: the exponent 54 is a certified convenience, not optimal, and does not beat the prime-field exponents of sources 04 and 17; Theorem B concerns one method class; large-torsion domains can make the hypothesis vacuous. Source 59: the radii $4/9$ and $1/2$ and the cutoff 16 are sufficient, not optimal; the global consequences follow from classical correction and source 11’s cubic; its Green–Tao remark concerns one displayed estimate; the lonely-runner bound is Bedert’s. All six: finite checks prove nothing general, no new bound for $r_k(N)$, no priority claim, no Lean proof. [7 October 2026; sources 60–69] *Further limitations of sources 60–69* (kept in place). Source 60: local theorems, no new global bound, no closing of the induction, no Lean formalization, no priority inferred; the compactness limits have no rate; B_s is not the envelope at every distance (equality only at $1/(2n)$); the activation threshold $1/(240(s-1))$ and the cubic $21/128$ at $s=4$ are not claimed optimal, nor P_s the full envelope; the restriction corollary is prime-uniform and its all-group form needs discrepancies of order above $L-1$; the Fejér scale is not proved optimal; finite checks support, not prove. Source 61: no proof-assistant verification or global bound; its exponent 911 is not competitive with 103 on prime moduli; keeping one scalar correlation is not keeping the derivative spectrum; no replacement of Sections 12–16; the lower bound on b_m does not show a_m optimal; the implementation is for small instances. Source 62: no end-to-end bound; the refined constants c_m are not explicit; the double logarithm is necessary only for its detector; the limit of $\kappa_{m,L}^*$ is not a uniform bound; the optimality of $\alpha^m \eta^{m-1}$ is not addressed; cyclic targets only; no bit-complexity bound. Source 63: no fixed- d value; the rate constants $1/8$ and 10 are not optimal; the local-scale limit precedes the degree limit and the fixed-degree optimum is open; no propagation through the later increments; its reviews are by ChatGPT instances, not peer review; no priority claim. Source 64: no improvement of the Szemerédi theorem or of graph Ramsey numbers; the disproofs of the conjectures are not its own; its exponents are not claimed optimal; $\beta_6 = 3$ and superpolynomial decay are open; the barrier concerns the specified model; no exhaustive priority search. Source 65: no proof-assistant verification; no priority; a characteristic-two local result, not a drop-in replacement for the prime-cyclic proposition; the cap $1 - 3d/2^{d+1}$ is not the optimal second value; maximizing functions are not classified. Source 66: no global bound or priority; its classifications are for ordered and torsion-free sets, not for subsets of cyclic groups; the endpoint hypothesis is not a total-energy hypothesis; the first gap is inverse Pollard theory. Source 67: “exact” only on the stated regions and interfaces; the recovery thresholds concern the four-point test; general odd targets and unrestricted quintics are open; further product-property information may compress the lists; one finite certificate component (eleven irreducibility checks); no global bound. Source 68: no global bound, no optimal finite restriction exponent, no priority; prime-field restrictions of sources 04 and 17 are better there; $\mathcal{I}_{m,a}$ measures only principal-spectrum cost; coordinate splitting is not resolved; the outer condition $\varepsilon < 1/4$ must stay strict. Source 69: no proof-assistant verification, unrestricted inverse theorem or global bound; symmetry is a hypothesis; the gap is not claimed to be the largest noncanonical energy; tensor classes, not functions, are classified.

V.3 Relationship to the ProveIt formalization

Source 01, delivered Section 14.

V.3.1 The inspected snapshot and normalization

The source inspection is pinned to commit

of VladimirReshetnikov/ProveIt, rather than to a moving branch [15]. The original paper was read through the public PDF mirror listed in the bibliography; selected repository files, not the entire repository build, were inspected.

The inspected `Definitions.lean` uses the unnormalized Fourier transform, a balanced indicator, finite partitions, natural-number progressions, and average cell size [16]. Consequently a hypothesis $|\widehat{\mathbf{1}_A}(r)| \geq \alpha N$ corresponds to $\eta = \alpha$ in this article, not to $\eta = \alpha/N$. A sum of absolute balanced cell sums becomes our D after division by N .

The companion `Proofs02DensityIncrement.lean` develops the positive/negative mass split for the Fourier density increment [17]. The companion `Proofs17FinitePatterns.lean` uses a finite sign-pattern/Sauer–Shelah approach for the inherited floor-pattern bounds [18]. These are natural locations for stronger reusable interfaces; no assertion that they have been replaced or recompiled is made.

[write] **6 October 2026.** At *HEAD* (25df8755d, 6 October 2026) the inspected files are unchanged except `Definitions.lean`, which gained a `Box.IsProper` that no source uses. The ledger now records 98 exact companions and 22 open statements, against 88 and 32 at this source’s pin. Of the statements compared here, `corollary_2_5`, `corollary_17_5`, `corollary_17_6` and `proposition_17_7` have exact companions, and `theorem_18_1` is open. Crosswalk: Section V.1.

V.3.2 Proposed formal interfaces

The dependencies can be organized in the following order.

Interface	Mathematical content	Reusable prerequisite
Real cell mass	$\sum g_j = 2 \sum (g_j)_+$ under $\sum g_j = 0$; weighted mean and occupied mass	Finite sums and order arithmetic
Cell phase envelope	Lemma I.1.2, first for an indicator, then for bounded weights	Complex norm, sine, finite convex combinations
Broken-chord transfer	Theorems I.2.1 and I.2.3 with explicit denominators	Convexity and mean identities
Integer partition	Theorem I.5.1, including the case of one block in a residue class	Finite pigeonhole and quotient/remainder
Fractional patterns	Positive perturbation and exact range factorization	Integer floor, finite minima of positive gaps
Pattern count	Hyperplane regions or an equivalent finite sign-family bound	Dimension and finite set cardinality
Translation quotient	Normalize a pattern by its empty-set value	Finite group action or an explicit canonical representative
Certificate checker	Rational primal/dual checks and exhaustive tree validation	Exact arithmetic and structural recursion

A practical first target is the $s = 0$ discrepancy theorem: its proof needs no trigonometry, no Fourier transform, and no structural additive combinatorics. The next target is the cancellation

inequality, followed by the full envelope. Proving the statements over arbitrary finite weighted families prevents repeated normalization and counting arguments later.

For the carry side, one should formalize interior realization before using a strict sign-pattern theorem. The signed-feature counterexample in Section IV.3 is a small regression test against a false boundary shortcut. The exact coefficient-range factorization also needs the open-endpoint surjectivity argument, not merely an upper-bound map.

For progression lengths, retain natural-number inequalities $N \geq L \prod_i Q_i$ and $L \leq \ell \leq 2L - 1$ as long as possible. Introduce square roots and floors only in the corollary. This isolates small-scale cases rather than trying to recover them from real asymptotic estimates. If actual code is added, it should prove a new strengthened statement and derive the old statement as a corollary where appropriate; changing a catalogue definition alone is not a proof.

V.3.3 A dependency warning, not a claimed global repair

The inspected `Sections17_18.lean` records the numbered statements as proposition-valued definitions. Its note attached to Theorem 18.1 flags an unsupported tower estimate in the published Corollary 16.11 as a quantitative dependency [19]. This is a report of the repository’s warning, not an independent conclusion here that a published headline theorem is false. Our local theorems do not use that estimate. In particular, applying Theorem I.4.1 to a supplied partition is unconditional, but asserting the existence of that partition with a given tower exponent would require a separate validated source.

No Lean or Lake executable was available in the working environment, and no new Lean theorem was compiled in producing this package. The document contains a formalization plan, conventional mathematical proofs, and separately described computational checks. It does not present a proposition-valued specification as a kernel-verified theorem.

[write] **6 October 2026.** Checked at HEAD: `theorem_18_1` is still open in the ledger, and its docstring in `Sections17_18.lean` still warns about the unsupported tower estimate in the published Corollary 16.11 (line 168). `FORMALIZATION_STATUS.txt` still lists resolving “the documented Corollary 16.11 bound” among its open obligations, and `corollary_16_11` is open. The dependency warning above is therefore accurate.

V.4 Further research questions and concrete next steps

Source 01, delivered Section 15.

The questions below are posed by this report. They are unresolved within this development; this is not a certification that every formulation is new or currently open in the entire literature.

Question V.4.1 (The second-order carry entropy). *Does $\log_2 C_k$ have a $-k \log_2 k$ correction, or is $\log_2 C_k = k^2 + O(k)$? More generally, determine the coefficient of $k \log k$ if it exists.*

Theorem IV.2.3 leaves exactly this scale between its upper and lower estimates. A promising route is to count intersections of the internal subset-sum hyperplanes that actually meet the open cube, rather than use the unrestricted region bound. The elementary anchor construction is too sparse to decide the leading coefficient, let alone this correction. Threshold-function embeddings give a strong baseline but need not capture all simultaneous carry levels.

Question V.4.2 (Exact counts and arrangement structure). *Can one obtain a recurrence, a useful intersection-poset formula, or a symmetry-reduced exact enumeration for C_k and A_k ?*

The supplied certificate format scales by adding more branches, but its size grows rapidly. Coordinate permutations and complement transformations suggest reductions; any quotient must retain enough information to certify that all orbit representatives are covered. An exact-arithmetic checker in a proof assistant would make such computations useful as formal regression tests, not merely numerical evidence.

Question V.4.3 (The modulus dependence of affine patterns). *For fixed k and coefficient range, when do modular collisions stabilize, and what is the exact number of translation classes as a function of M ?*

Theorem IV.3.2 is uniform in M but deliberately discards these collisions. A finite list of integer-valued normalized patterns exists; modular coincidence is then controlled by divisibility of finitely many coordinate differences. Determining those difference vectors could give an exact piecewise description, including a threshold beyond which no new collisions occur. The offset's fractional part must remain visible.

Question V.4.4 (Can the selection loss be reduced?). *What additional orthogonality or positivity assumptions allow a phase-removal interface to lose only L^{-1} rather than L^{-2} in energy, without changing the desired structural conclusion?*

For arbitrary complex summands the squared pigeonhole loss in Proposition IV.3.3 is sharp: equal aligned summands attain it. A stronger theorem therefore needs additional hypotheses, not just more careful algebra. One possible approach is to retain the sum of squared contributions across translation classes until after a structural selection step. Such a mechanism has not been established here.

Question V.4.5 (Sharpness under arithmetic phase constraints). *How much stronger can the density-tail bound become when ψ is one fixed cyclic character and the cells are proper progressions with a common step and nearly equal lengths?*

The extremizers of Theorem I.2.1 freely rotate cell arcs and assign endpoint phases to occupied points. A genuine character imposes spacing and compatibility restrictions. A quantitative analysis of those restrictions could improve constants or density dependence beyond the abstract sharp bound; it would not contradict that bound's optimality in the larger weighted model.

Question V.4.6 (The optimal finite size-increment tradeoff). *Can one combine the exact density tradeoff (I.3.4) with an optimal finite partition algorithm, including the final incomplete blocks, to characterize the best uniform pair of progression length and increment?*

Our integer partition is transparent but not claimed to have optimal constants. Theorem I.7.2 fixes the square-root exponent for one-frequency phase-flat methods. Remaining improvements may come from choosing unequal block lengths, exploiting rational frequencies, or paying for selected short cells through their occupied mass rather than their ambient size.

Question V.4.7 (Retaining more than one density statistic). *Which finite collection of stop-loss moments most efficiently preserves useful density-tail information through repeated localization steps?*

Theorem I.4.2 supplies an exact optimization method once the statistics are chosen. The research issue is to prove structural lemmas that propagate those statistics at acceptable cost. A two- or three-knot state could be substantially more informative than a single discrepancy number, but a global gain must include the cost of obtaining that state.

Question V.4.8 (Higher-degree phase partitions). *Can polynomial or nilsequence localization furnish partitions with both quantitative minimum cell sizes and small occupied exceptional mass, so that the sharp tail theorem applies without a large selection loss?*

The density transfer itself is already proved for any short-arc phase; what is missing is the required structural partition. Simultaneous linear approximation alone does not solve the higher-degree problem. The polynomial floor-pattern bound in Theorem IV.4.1 controls one finite complexity term but does not construct those partitions.

Question V.4.9 (Polynomial floor entropy beyond the leading term). *For fixed degree d , determine the secondary terms of $\log_2 |\mathcal{P}_{n,d}(r)|$, and identify the transition when $\log r$ is comparable to n .*

The finite upper bound separates a coefficient-range term from a region count, but unlike Theorem IV.1.2 it need not be an exact product. Understanding overlaps between different integer coefficient vectors is essential. The threshold-function lower bound alone does not settle the large-range contribution.

Question V.4.10 (A fully audited quantitative integration). *After independently checking the earlier structural parameter estimates, what explicit global improvement survives if the new density-transfer and translation-class bounds are propagated through one complete argument?*

The first deliverable should be a dependency graph with normalized quantities and verified recurrences, including the issue flagged at Corollary 16.11 in the inspected repository. Only then is it meaningful to compare a resulting global bound with modern benchmarks. This prevents a smaller local factor from being mistaken for a new global theorem.

V.5 Conclusions and result-status ledger

Source 01, delivered Section 16.

The work yields a reusable sharp transfer principle and a more economical floor-pattern count. The density calculation is sensitive to the actual balanced mass $2\delta(1 - \delta)$, to the full distribution of cell densities, and to occupied exceptional mass. The counting calculation separates integral coefficient freedom from fractional carries, treats boundary points correctly, and exhibits a threshold-complexity obstruction to a subquadratic exponent. Both refinements have explicit hypotheses and can be inserted at well-defined proof interfaces.

Claim	Status in this package
Sharp finite weighted density-tail theorem, including equality examples	Proved in Theorem I.2.1; arithmetic extremizers are not claimed.
Occupied-mass exceptional-set bound and optimal forced maximum	Proved in Theorems I.2.3 and I.3.1.
Explicit Fourier density increment with integer-exact length	Proved in Theorem I.6.1, for arbitrary real frequency and bounded weights; no prime-modulus restriction.
Exact range factorization and quadratic-exponent carry upper/lower bounds	Proved by elementary arguments in Theorems IV.1.2, IV.1.4, and IV.2.1.

Claim	Status in this package
$\log_2 C_k \sim k^2$ and polynomial-floor leading entropy	Deduced from explicitly identified classical threshold-function inputs; those inputs are not reproved here.
$C_4 = 154$	Exhaustive exact rational certificates with a Python verifier and a completeness argument; not Lean-verified.
Smaller translation-class budget	Proved in Theorem IV.3.2; subsequent global parameter propagation is not supplied.
Improved record bound for Szemerédi’s theorem	Not claimed.
Priority of every formulation	Not certified by an exhaustive literature search.

V.6 Boundary and normalization checklist

Source 01, delivered Appendix A.

The formulas in this article are intended to be used with the following conditions kept explicit.

1. The phase results assume $0 < \delta < 1$. At $\delta = 0$ or 1 , a bounded function with that mean is constant and the balanced correlation is zero; formulas dividing by $\delta(1 - \delta)$ should not be used.
2. The short-arc theorem uses half-angle $\theta < \pi/2$. The progression partition uses full arc length r in turns, so $\theta = \pi r$, not $2\pi r$.
3. A normalized Fourier coefficient is the unnormalized coefficient divided by N . The normalized discrepancy is likewise the sum of cell balanced masses divided by N .
4. The density-tail event is $q > a$, strictly. Its stop-loss lower bound may vanish at the optimal forced maximum even though a cell with $q \geq a$ is guaranteed.
5. Exceptional occupied mass is $\mathbb{E}(h\mathbf{1}_E)$ in the original probability measure. It is not a conditional density on E .
6. The coefficient-range factorization assumes integral $r_i \geq 1$ and open intervals $(-r_i, r_i)$. Fractional coordinates can lie on a boundary only because Lemma IV.1.1 removes that boundary without changing the pattern.
7. Natural progressions have positive step and no repeated terms. Singletons handle the small-scale case; no nonintegral number is treated as a possible cell length.
8. The affine translation quotient retains the fractional offset. Its modulus-free upper bound does not assert absence of modular collisions.
9. The threshold-function asymptotic is an external theorem. The elementary finite pattern bounds and all density-transfer statements remain valid without it.

V.7 Source and artifact provenance

Source 01, delivered Appendix B.

The repository snapshot was read on 6 October 2026. The paper directory contains both a PDF and a TeX representation. The research arguments use the original paper’s relevant PDF pages and selected Lean source files, not an assumption that every statement in the repository has already been proved. The pinned commit identifies precisely which source comments and normalizations were inspected. No repository files were modified.

The package includes the article source, compiled PDF, certificate data, verification and supplementary-test scripts, result logs, and a README. The article’s source uses an internal bibliography and does not require BibTeX or downloaded repository files to compile. No copies of the copyrighted source papers or font files are redistributed.

V.8 Formalization interface and reproducibility

Source 02, delivered Section 10.

V.8.1 What was inspected

The repository’s `Definitions.lean` defines `fourier` using the unnormalized `ZMod.dft`, defines `difference` with $x - h$, and defines `balanced`, `IsPartition`, `UniformOfDegree`, and the unnormalized `gowersNorm`. The inspected `ProofInfrastructure.lean` provides, among other helpers, `IsPartition.sum_card` and `sum_balanced`. The inspected part of `Proofs02DensityIncrement.lean` contains private helpers for partition sums and the exact equality between total absolute mass and twice positive mass [20].

The normalized U^3 quantity in this article translates to the repository’s degree-two uniformity numerator by

$$\mathcal{U}(f) = N^{-4} \sum_{a,b \in G} \left| \sum_{x \in G} \Delta_a \Delta_b f(x) \right|^2.$$

Thus the degree parameter of `UniformOfDegree` is 2 here, not 3. For the repository’s unnormalized order-three Gowers norm, the eighth power is $N^4 \mathcal{U}(f)$. These normalizations should be established as explicit bridge lemmas before the energy theorem is imported.

V.8.2 Suggested dependency order

The most economical formalization starts with finite weighted real sums: positive/negative mass, the negative-support bound, and Theorem I.9.1. It then proves the real-phase lemma using positive and negative weighted averages. These results have no Fourier dependencies and can be added independently of a full higher-order inverse theorem.

For Theorem II.2.1, first formalize the three identities (II.2.2)–(II.2.4) over `ZMod N`. Keep all statements polynomial and multiplication-based, rather than dividing by $m(f)$ or $\mathcal{U}(f)$; this handles zero cases without extra hypotheses. A later generalization can replace `ZMod N` by an abstract finite abelian group and its character group.

Next prove the two-tetrahedron bijections in Lemma II.3.1, then add graph indicators and balanced indicators as specializations. The final arithmetic-progression wrapper should accept a proved $\text{PP}(K, R_0)$ hypothesis rather than hiding it behind an unverified claim about exact cell counts.

Care is needed with empty cells: the repository’s general partition predicate does not itself require every cell to be nonempty. Empty cells must be removed before using cell means, $M - 1$, or a minimum cell length. Likewise, injectivity of a modular progression parameterization is not the same as nonwrapping in its standard integer representatives.

V.8.3 Verification status

No Lean compiler or Lake executable was available in the working environment. Accordingly, none of the new results is represented as kernel-checked, and no untested Lean proof is included. The source reading was targeted rather than a certification of every theorem in the supplied development. The mathematical proofs in this article are the primary deliverable.

The accompanying `verify.py` independently checks the weighted Fourier identities and both asymmetric forms of the energy inequality on finite cyclic groups, compares spectral and direct convolution energy for small groups, tests phase and tail inequalities, and evaluates exact sharpness examples. It uses the fixed seed 20261006. The run recorded in `verification_results.json` includes 490 random energy cases, 2,500 random phase cases, 3,000 zero-mean moment cases, 1,500 nonzero-mean moment cases, and 1,500 phase-budget optimization cases, together with three exact-form equality families evaluated numerically and a rational moment equality example. These experiments are regression checks, not proofs or substitutes for formal verification.

[write] **6 October 2026.** The suite is shipped as `code/02-energy-freezing-verify.py`; it overwrites `verification_results.json` beside itself, so the README’s rerun route runs it on a copy. At intake it reproduced the recorded run identically (8,990 random cases and the exact families). The corrected Corollary 5.6 that Section I.11 imports is now formalized (note there).

V.9 Further research questions and proposed work

Source 02, delivered Section 11.

The following are proposed research targets suggested by the proven interfaces. Their inclusion is not a claim that each is a previously published open problem.

V.9.1 A correlation-preserving structural theorem for multivalued spectra

Proposition II.4.1 retains all significant derivative frequencies. Can the subsequent energy-to-structure and graph-selection steps preserve a fixed fraction of t_τ while exploiting the denominator κ ? The decisive issue is not merely obtaining a small-doubling subset: it is ensuring that the subset continues to carry the correlation needed for integration. A useful first goal is a theorem whose hypotheses include both additive energy and a specified nonnegative correlation weight, and whose conclusion controls both losses explicitly.

V.9.2 Classify near equality for the coupled pair (f, w)

Proposition II.4.2 gives two quantitative alignments, while Examples II.3.3–II.3.4 show distinct sharp geometries. Existing results about near-maximal Gowers norms [7] are relevant but do not by themselves state a classification of the additional spectral weight in our mixed inequality. One can ask for a dimension-free theorem describing w and f jointly when the main inequality is nearly saturated, with appropriate treatment of subgroups and quadratic phases.

V.9.3 Retain norm profiles through higher-order integration

The graph-energy step is repeatedly applied to derivatives in higher-degree arguments. A natural target is a weighted family theorem that keeps the distribution of $m(\Delta_{\mathbf{h}}f)$ and $\mathcal{U}(\Delta_{\mathbf{h}}f)$ instead of replacing every slice by a supremum. The immediate difficulty is compatibility: good structural pieces from different slices must agree on overlapping configurations. A local improvement cannot be counted as a global gain until this compatibility cost is controlled.

V.9.4 A complete quantitative ledger for one progression length

An attainable next project is to propagate the present denominators, the sharp real-phase law, and the optimized budget through every step of the length-four proof, with one explicit choice of parameters. This would determine whether a real improvement survives the remaining inverse and recurrence estimates. Only after that full ledger is complete would it be justified to state an improved numerical threshold for that proof. Comparisons should distinguish improved constants within Gowers’s route from stronger bounds obtained by later methods [4, 2].

V.9.5 Sharper recurrence inputs and heterogeneous budgets

Theorem I.11.1 treats PP as a replaceable component. Any better value of K or R_0 transfers automatically to the progression certificate. A separate optimization problem assigns different diameter budgets to cells of different sizes and correlations, using the heterogeneous inequality in Section I.8. The local quadratic bound contains still more information than its linear envelope; exploiting it may improve nonuniform cases without altering the recurrence theorem.

V.9.6 The finite-cell extremal problem

Theorem I.9.1 solves the weighted total-good-mass problem sharply. It does not solve the problem of the largest good cell when M and all allowed cell sizes are fixed. The three-type extremizer explains why a negative cell and a low-positive cell can consume separate parts of the cell budget. A precise finite-dimensional optimization, including integer cardinalities, could strengthen the $M - 1$ corollary and produce exact finite- N certificates.

[write] **6 October 2026; sources 06–10.** *Advanced by source 06.* With M fixed and the cell masses free (the weighted model), the problem is solved by source 06’s Theorem I.21.1: the exact best lower bound for the largest high cell is $\min\{P_+/(b(M - 1)), S/(M - 2)\}$, attained by explicit atomic extremizers; its second extremizer is a three-type configuration of exactly the kind described above (cells at b , one cell at the threshold, one negative cell). Source 06 realizes its extremizers on finite uniform spaces when the parameters are rational. Still open: prescribed allowed cell sizes and integer cardinalities at a fixed N (Section V.2, item I.6).

V.9.7 A characteristic-independent Lean library

The main analytic results work on every finite abelian group and the two-tetrahedron argument uses no field division. A useful library contribution would formalize weighted Fourier energy, its normalization on $G \times \widehat{G}$, and real-valued phase freezing independently of the prime-cyclic presentation. Regression examples should include even cyclic groups, subgroup spectral rectangles, empty weights, zero functions, and wrapping versus nonwrapping progression charts.

V.10 Conclusion

Source 02, delivered Section 12.

The strongest unconditional analytic conclusion of this article is the mixed weighted inequality (II.2.1), with its exact norm denominators and coefficient-one equality examples. The strongest elementary density conclusion is the sharp tail bound (I.9.1), coupled with the sharp aggregate real-phase law (I.8.3). Together they strengthen identifiable interfaces in Gowers’s argument, including Proposition 6.1 and Lemma 5.15, and yield an explicitly optimized progression certificate under a clearly stated partition input.

The remaining challenge is structural, not merely algebraic: local gains must survive extraction, integration, recurrence, and iteration. The article makes those gains available without presenting that uncompleted propagation as a new global theorem.

V.11 Source audit and normalization checklist

Source 02, delivered Appendix A.

The repository branch reference observed during this session was

66c02f8a1d26f553deac5a91dada1739e0d15064.

The specific inspected blobs are recorded below, so the source comparison is not dependent only on a moving branch name.

File below the supplied repository paths	Git blob SHA
sz-thm-gowers-proof.tex	41085a0efde8c1932e86e80c791984f52841c22e
Definitions.lean	202260e0d7ec6f033c1758347381c34a6f0d9212
ProofInfrastructure.lean	ebd9e7ed02d81de2e17b97b4223f861b785ed5d8
Proofs02DensityIncrement.lean	85b5c3ad8769b9769e1939ae5c05f0757f0678b5

[write] **6 October 2026.** Checked at HEAD (25df8755d): the four blob identifiers are those of commit 66c02f8a1. The transcription, `ProofInfrastructure.lean` and `Proofs02DensityIncrement.lean` are unchanged at HEAD; `Definitions.lean` has changed (blob 97113b7), by an added `Box.IsProper` that source 02 does not use. Source 02’s delivery README records no pin; this branch reference is its provenance (Section F.4).

The targeted paper passages were the corrected Corollary 5.6 and nearby phase-partition statements, Lemma 5.15, and Proposition 6.1. The transcription includes editorial changes; a theorem imported from that edited source is labeled as such above rather than silently attributed to the exact original printed formulation.

Quantity	This article	Paper/Lean convention
Fourier coefficient	$\mathbb{E}_x f(x) \overline{\chi(x)}$	$N \widehat{f}(\chi)$
Derivative-graph statistic	T	$N^3 T$
Order-three cube quantity	$\mathcal{U}(f)$	$N^4 \mathcal{U}(f)$
Graph additive energy	$\mathcal{E}(\mathbf{1}_\Gamma)$	$N^3 \mathcal{E}(\mathbf{1}_\Gamma)$
Balanced L^2 mass	$m(f) = \delta(1 - \delta)$	$Nm(f)$ before averaging

V.12 How local improvements affect an iteration

Source 03, delivered Section 9.

The estimates above remove real losses, but they occur at different places in the proof. The cube theorem improves a counting consequence of uniformity. The DRC theorem improves extraction of a subfamily from intersection data. The Fourier theorem converts an already established correlation into a density increment on an ordinary progression. Their parameters are not interchangeable. In particular, a larger common-neighborhood core does not by itself provide an improved correlation theorem for arbitrary high Gowers norms.

Figure 1 displays two of the local comparisons using the literal formulas proved in the text. For cube counts, retaining the centering changes the first possible error power, while classifying its supports further reduces its coefficient. For DRC, retaining the original success criterion and changing the exponent from five to three gives a growing advantage as the density parameter decreases.




Figure 1: Evaluations of proved local bounds. Left: relative cube-count error bounds for $d = 3$, $\delta = 1/2$, comparing the triangle inequality, the cancellation of degrees one through three, and the exact quartic types on an odd-order group. Right: guaranteed core fractions at the same threshold $\delta^2 m/2$ and the same 90% good-pair requirement. These are upper or lower bounds as labelled, not measurements of extremal examples or claims that each curve is attainable.

V.12.1 An explicit abstract recurrence

The following elementary calculation explains which parameters must be tracked before any global conclusion can be stated. It is a conditional recurrence theorem, not a substitute for the analytic density increment theorem needed to supply its hypotheses.

Proposition V.12.1 (Density and length bookkeeping). *Let $c > 0$, $p > 1$, $C > 0$, $q \geq 0$, and $0 < \beta < 1$. Suppose positive numbers $\delta_i \leq 1$, $N_i > 0$ satisfy, for $0 \leq i < R$,*

$$\delta_{i+1} \geq \delta_i + c\delta_i^p, \quad N_{i+1} \geq C\delta_i^q N_i^\beta.$$

Then

$$R \leq \frac{(1+c)^p}{c(p-1)} (\delta_0^{1-p} - 1). \quad (\text{V.12.1})$$

Writing $B = \max\{0, -\log(C\delta_0^q)\}$, one also has

$$\log N_r \geq \beta^r \log N_0 - B \frac{1 - \beta^r}{1 - \beta} \quad (0 \leq r \leq R). \quad (\text{V.12.2})$$

Consequently, for an integer $T \geq R$ and $H \geq 1$, the condition

$$\log N_0 \geq \beta^{-T} \left(\log H + \frac{B}{1 - \beta} \right) \quad (\text{V.12.3})$$

ensures $N_r \geq H$ at every $r \leq R$.

Proof. The potential u^{1-p} decreases at a controlled rate. Since $\delta_i \leq 1$,

$$\begin{aligned}\delta_i^{1-p} - \delta_{i+1}^{1-p} &\geq (p-1) \int_{\delta_i}^{\delta_i + c\delta_i^p} t^{-p} dt \\ &\geq \frac{(p-1)c\delta_i^p}{(\delta_i + c\delta_i^p)^p} \geq \frac{(p-1)c}{(1+c)^p}.\end{aligned}$$

Summing and using $\delta_R^{1-p} \geq 1$ gives (V.12.1). The densities are increasing, so $\log N_{i+1} \geq \beta \log N_i - B$. Iterating this affine inequality gives (V.12.2). Under (V.12.3), $\beta^r \log N_0$ is at least $\log H + B/(1-\beta)$ for $r \leq T$, which proves the last assertion. $\square$

For fixed constants, the familiar double-exponential length budget is visible in (V.12.3): the number of iterations occurs in the exponent of β^{-T} , which already bounds a logarithm of the starting size. Improving C or q changes the additive length cost B . Improving the density growth exponent p can change the scale of the number of iterations. These are different effects. To extract a better global Szemerédi bound from this report one would have to identify, and prove throughout the later sections, the corresponding updated values or functions of all these parameters. We have not carried out that full replacement of the higher-order inverse and localization arguments.

[write] **6 October 2026; sources 06–10.** *Two further iteration statements.* Source 10’s Proposition V.42.1 bounds the number of steps of the same recurrence by $(\delta_0^{1-p} - 1)/(1 - (1+c)^{1-p})$, and source 06’s Lemma V.25.1 with $G = 1$ by $1 + (\delta_0^{1-p} - 1)/(c(p-1)) + p \log(1/\delta_0)$ (its (V.25.3)). *Comparison (second write).* With $q = p-1$, $1 - (1+c)^{-q} = q \int_0^c (1+u)^{-q-1} du$ lies between $qc(1+c)^{-q-1}$ and qc , so the coefficients of δ_0^{1-p} satisfy $1/(cq) \leq 1/(1 - (1+c)^{-q}) \leq (1+c)^p/(cq)$: source 06’s leading coefficient is the smallest and source 10’s lies between it and the one in (V.12.1); source 06’s bound has the additional term $1 + p \log(1/\delta_0)$ and counts the densities $\delta_0, \dots, \delta_{t-1}$, whereas the bound here counts steps. For $p = 2$, $c = 1/320$ the coefficients are 320, 321 and $321^2/320 \approx 322.003$. Source 06 also proves that the leading term of its double-logarithmic length budget is optimal for its abstract local rule (its Section V.25), and source 10’s Proposition V.42.2 is the length recursion (V.12.2) with variable parameters.

The same distinction applies to the exact quartic identity. It is stronger information about cubes with a prescribed balanced Gowers norm. It does not imply that every set lacking arithmetic progressions has a large ordinary Fourier coefficient. Gowers’s motivating examples address precisely why higher-order structure is needed. The role of the improved count is to make the uniformity branch and its constants more precise.

V.13 The Lean development and a concrete formalization route

Source 03, delivered Section 10.

V.13.1 Pinned source and status of the inspection

All repository comparisons refer to revision

3a25020a654d196f7f1aa2e37649752c000d0e65

of ProveIt, inspected on 6 October 2026. The requested paper is at `Combinatorics/Ramsey/Papers/sz-thm-gowers-proof`; the Lean files are at `Combinatorics/Ramsey/Lean/GowersSzemeredi`. The adjacent source inventory, `gowers-proof-status.json`, reports 88 numbered statements with exact companion declarations and 32 open statements at that revision [22].

[write] 6 October 2026. *Stale since the pin.* At HEAD (25df8755d, 6 October 2026) the inventory records 98 exact companions and 22 open statements. The ten statements closed since 3a25020a6 are `corollary_5_6`, `corollary_5_7`, `lemma_5_9`, `lemma_5_10`, `corollary_5_11`, `lemma_5_14`, `lemma_16_1`, `lemma_16_3`, `lemma_16_5` and `lemma_16_6` (compared at the write). `theorem_1_3`, `theorem_18_1` and `theorem_18_2` are still open, so the conclusion of the next paragraph stands. Crosswalk: Section V.1.

Those counts are a source inventory, not a certification of the whole development. The facade exposes numbered results as propositions, usually with companion theorems whose names end in `_holds`. Defining a proposition neither proves it nor adds it as an axiom. The inventory explicitly separates exact companions from conditional, restricted, or otherwise related results. In particular, the quantitative headline `theorem_1_3` and final statements `theorem_18_1` and `theorem_18_2` are marked open there. This inspection therefore does not support describing the repository as a completed formal proof of the full quantitative theorem.

The files directly relevant to our comparisons contain substantive companion proofs. We inspected their definitions and proof scripts; we did not compile the Lean project. The new mathematical statements in this report have written proofs and independent mathematical reviews, together with the finite checks described in the appendix. No new machine-checked Lean proof is claimed or supplied.

[write] 6 October 2026. The “independent mathematical reviews” mentioned here and in Section V.15 are not part of the delivery and are not recorded anywhere in the repository; this report relies on the printed proofs, the intake’s checks and the write’s notes, not on them.

V.13.2 Normalization map

The following conversions prevent a common source of false improvements. Let $N = |G|$, and take $d \geq 1$.

Quantity	This report	Inspected repository convention
Fourier transform	$\mathbb{E}_x f(x) \overline{\chi(x)}$	Unnormalized sum, equal to N times this transform
Gowers norm	$\ f\ _{U^d}$	<code>gowersNorm</code> d $f = N^{(d+1)/2^d} \ f\ _{U^d}$
Uniformity parameter	$\ f_A\ _{U^d}^{2^d} \leq \alpha$	<code>UniformSetOfDegree</code> A α $(d-1)$
Cube count	$Q_d(\mathbf{1}_A)$	<code>cubeCount</code> $= N^{d+1} Q_d(\mathbf{1}_A)$
Cube arguments	$x + \sum_i v_i h_i$	A minus convention in <code>cubeArgument</code> ; negating all h_i gives the same average

Both cube conventions count all parameter tuples, including degenerate cubes. This is why the rounding lemma can bound the exceptional tuples by a collision probability instead of changing the definition of the norm. Changing to unlabelled, injective, or symmetry-quotiented cubes would require an additional counting theorem.

V.13.3 Useful existing declarations

The source provides the following starting points in namespace `LeanProofs.GowersSzemerédi`. These are concrete interfaces to reuse or generalize, rather than assertions that the proposed additions have already been formalized.

Partition discrepancy and Fourier increment.

`Proofs02DensityIncrement.lean` proves `corollary_2_5_holds`. Its private `cor25_large_scale` already obtains the stronger $\rho/4$ increment and $\sqrt{\rho N/(128\pi)}$ length before weakening its conclusion. General finite weighted sums, followed by Theorem I.14.1, would expose the asymmetric information without repeating the Fourier argument.

Multilinear norms and cube expansion.

`Proofs03GowersCS.lean` contains the public homogeneous theorem `gowers_cauchy_schwarz`, which already allows functions beyond the unit disc. `Proofs03Minkowski.lean` contains the triangle inequality and private distributive expansion helpers. `Proofs03CubeUpper.lean` proves `lemma_3_10_holds` from that triangle inequality. Exporting a general expansion indexed by vertex subsets is a natural first step toward the exact quartic theorem.

Selection from many cells.

`Proofs05_10.lean` contains `lemma_5_15_holds`. Its assumptions and conclusion match the symmetric discrepancy extraction discussed in Section I.14. The general theorem should be stated with real nonnegative cell weights and asymmetric cell-sum bounds, then specialized to cardinalities.

Dependent random choice.

`Proofs07DRC.lean` proves `lemma_7_4_holds` using ordered samples with replacement, indexed by `Fin5`. Its internal score bound already has constant $11/16$ multiplying $\delta^{10}n^2$. The main new ingredient for the cubic theorem is a pointwise polynomial minorant, followed by the exact three-sample counting identity. Merely exposing the internal five-sample constant would give a smaller change.

V.13.4 Recommended order of implementation

First formalize the finite weighted selection theorem over an arbitrary finite index type. Prove positivity of its denominators once, and keep the pointwise inequalities separate from the division steps. This produces immediately reusable specializations to Lemma 5.15 and balanced indicators.

Second formalize the cubic certificate. Its good-region proof is a polynomial identity followed by nonnegativity; its bad-region proof may use concavity or a direct monotonicity argument. Keep the sampling identity generic in the number of witnesses. Specialization to three witnesses then replaces only the numerical part of the current DRC proof.

Third export the general cube expansion and prove the single-coordinate cancellation lemma. Nonempty supports of at most three selected vertices, or four vertices with nonzero XOR, leave one variable occurring in exactly one centered factor after a coordinate reflection. The corresponding reindexing is a bijection and uses no division by two. The first quartic upper bound then follows directly from Gowers–Cauchy–Schwarz, before any Fourier classification is needed.

Fourth formalize the two nonzero four-vertex types. One may begin over $\mathbb{Z}/N\mathbb{Z}$ to match existing files, then generalize to finite abelian groups. The extra quotient energy requires explicit handling of order-two characters; assuming an invertible doubling map is appropriate only for the odd-order specialization. The combinatorial formulas for R_d, T_d are separate finite enumeration theorems.

Finally integrate the progression-count telescope, the proper progression partition, and the four-edge compression statement. Each has a clear finite interface. The generalized von Neumann theorem should state invertibility of slope differences explicitly; the progression partition should

carry its lower and upper integer lengths in its output. This order establishes useful lemmas early while avoiding dependence on the unfinished quantitative headline.

V.14 Further research questions and topics

Source 03, delivered Section 11.

The following questions are proposed continuations of the proved results. Some are deliberately scoped finite problems, while others concern the larger proof architecture. A question proposed here is not being presented as a previously published open problem unless a source is identified.

Research question V.14.1 (The odd-order leading constant). For fixed $d \geq 3$, determine

$$c_d^{\text{odd}} = \sup_{\substack{G \text{ finite abelian, } |G| \text{ odd} \\ 0 \neq f: G \rightarrow \mathbb{R}, \mathbb{E}f=0}} \frac{\|f\|_{U^2}^4}{\|f\|_{U^d}^4}.$$

The exact quartic theorem reduces the best leading coefficient on odd-order groups to $R_d c_d^{\text{odd}}$. Is the supremum attained or approached by a real part of a character, by several Fourier modes, or by a different construction?

[write] **6 October 2026; sources 28–33.** *Advanced for $d \geq 4$, not settled* (sources 28 and 29, Section III.66). Source 29 proves $\|f\|_{U^d}^{2^d} \geq 2^{-m} \binom{2m}{m} \|f\|_{U^2}^{2^d}$, $m = 2^{d-2}$, hence $c_4^{\text{odd}} \leq (8/35)^{1/4}$ and $c_d^{\text{odd}} \leq B_d \downarrow 1/2$, and $c_d^{\text{odd}} \geq 1/3$ for every d (Corollary III.69.1, Proposition III.69.2); source 28 proves $c_4^{\text{odd}} \geq 0.5587294$ and that a real part of a character is not extremal for $d \geq 4$ (Corollary III.70.8, Theorem III.70.9); the intake’s certificate gives $c_4^{\text{odd}} \geq 0.5587725$. So $0.5587725 \leq c_4^{\text{odd}} \leq (8/35)^{1/4} = 0.6914416$ among sources 01–33; later placed sources record better upper bounds (item III.1).

[write] **6 October 2026; sources 34–39.** *The upper end improved again* (sources 34 and 35, Sections III.71–III.72). Source 34 proves $\|f\|_{U^2}^4 \leq (2/3) \|f\|_{U^4}^4$, that is the coefficient 81/16 in source 29’s normalization (Theorem III.71.1); source 35, the same day, proves the coefficient 5 (Theorem III.72.1); for $5 \leq d \leq 8$ source 34’s coupled bounds slightly improve source 29’s (Remark III.71.10). So $0.5587725 \leq c_4^{\text{odd}} \leq 2/3$ among sources 01–39; source 45, placed and not yet written, records $(832/4795)^{1/4} = 0.6454070$ (item III.1).

[write] **7 October 2026; sources 40–46.** *The upper end improved again* (source 45, Section III.89): $\|f\|_{U^4}^{16} \geq (4795/832) \|f\|_{U^2}^{16}$ (Theorem III.90.1), so $0.5587725 \leq c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$ among sources 01–46; source 49, placed by f8bc5e2ec and not yet written in, records $29/4$, that is $c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ (for its write) (item III.1).

[write] **7 October 2026; sources 47–53.** *The upper end improved again* (source 49, Section III.92): $\|f\|_{U^4}^{16} \geq (29/4) \|f\|_{U^2}^{16}$ (Theorem III.93.1), so $0.5587725 \leq c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ among sources 01–53, and $c_5^{\text{odd}} \leq 29/50$ (item III.1).

The universal bound $c_d^{\text{odd}} \leq 1$ follows from monotonicity. Our cosine example gives a positive explicit lower bound. The order-two character that attains the unrestricted coefficient is unavailable here, so the sharp universal result does not settle this restricted problem. A practical first project is a systematic search over small odd cyclic groups, followed by an inequality explaining any apparent extremizer. Numerical optimization would be evidence, not a proof of the supremum.

[write] **6 October 2026.** *Settled for $d = 3$ (cross-source finding).* Source 04’s Theorem III.3.4 gives $c_3^{\text{odd}} = 2^{-1/2}$, attained exactly by the real parts $2 \operatorname{Re}(c\xi)$ of nontrivial characters (the question’s first alternative); proof in the note after that theorem. The best odd-order leading coefficient for $d = 3$ is $R_3 c_3^{\text{odd}} = 6\sqrt{2}$, as source 04 proves directly. For $d \geq 4$ the question is open; it coincides with source 04’s Question 3 (Section V.2, item III.1).

Research question V.14.2 (Classify the fifth and sixth centered coefficients). For each $r = 5, 6$, classify the supports $S \subseteq \{0, 1\}^d$ of size r for which the character system (III.2.27) can have solutions with no trivial character. Derive multiplicities as functions of d and identify the necessary quotient correlations and Fourier phase invariants.

The coordinate-reflection test removes every support with a coordinate occupied by exactly one selected vertex or by all but one. For surviving supports, Smith invariant factors identify possible small-torsion obstructions. A successful classification should distinguish actual vanishing of a centered correlation from mere rational independence of its defining forms. The fifth-order examples discussed at the end of Section III.2 give concrete starting cases.

[write] **6 October 2026; sources 06–10.** *Fifth order answered, sixth order open.* For $r = 5$ the question is answered by source 09's Theorems III.20.1 and III.21.1 with Proposition III.20.2: two support types, multiplicities Q_d and S_d , and the correlations M_5 (order-two coupling χ, χ^2) and $T_{3,5}$ (order-three characters). For $r = 6$ it remains open; source 09 shows that an isolated order-five term occurs there (Theorem III.23.1) and asks the same question (Section V.2, item III.2).

Research question V.14.3 (The sharp fixed-threshold DRC profile). Let $F(\delta)$ be the largest guaranteed fraction $|K|/n$ under the exact hypotheses and success criterion of Lemma 7.4. Determine the correct order of $F(\delta)$ as $\delta \downarrow 0$. Can the exponent 3 in Theorem II.5.2 be improved using constraints specific to intersection matrices?

The scalar optimization proves only that a particular two-witness first-moment certificate cannot give a positive uniform estimate at the original parameters. Actual overlap matrices are symmetric positive semidefinite Gram matrices of indicators and obey further incidence constraints. It is a real additional problem to determine whether those constraints rule out the scalar bad distributions. Conversely, a lower bound on the necessary exponent requires realizable families of sets, not just a distribution of formal overlap values.

Research question V.14.4 (Optimize the entire elementary extraction). Optimize the three parameters $\gamma, \varepsilon, \beta$ in Proposition II.6.2, and compare alternative walk lengths and popularity thresholds. Which size–difference-set tradeoffs can be achieved while retaining only elementary finite inequalities?

The meaningful objective is the pair of conclusions, including the normalization of the difference-set bound by $|A|$ or by $|A'|$. Optimizing one displayed exponent while allowing the selected set to shrink excessively can conceal a loss. Comparison with [12] should remain explicit; the purpose of this project would be a simple formalizable proof or a better constrained tradeoff, not to mistake an elementary refinement for an improvement of the strongest known theorem.

Research question V.14.5 (Allocate phase oscillation by local mass). Suppose a phase partition can be refined at a cost depending on the desired oscillation in each cell. Can one minimize the total refinement cost subject to

$$\sum_j \int_{P_j} |f| |\psi_j - z_j| d\mu \leq \lambda \rho$$

by using local information on $|f|$, rather than a single global oscillation tolerance?

Lemma I.14.3 already provides the exact analytic constraint. The unresolved work is geometric: prove a refinement theorem with a cost function sharp enough for this optimization to matter. One can start with linear phases and known cell masses, then study polynomial phases. The global density factor in Theorem I.15.2 is the simplest instance of retaining this information.

Research question V.14.6 (Local counting that rectifies automatically). Develop a version of Theorem III.1.2 for interval-supported functions or localized norms in which the conclusion directly counts ordinary integer progressions and includes all boundary errors.

The elementary second-difference test gives a clean sufficient condition for rectification, but the global modular norm does not enforce that condition. The target is an explicit statement whose hypotheses reflect both uniformity and support. Any proposed bound should be tested on sets concentrated near the ends of the interval and on small composite moduli before it is fed into the density increment iteration.

Research question V.14.7 (Propagate improvements through a verified dependency graph). Replace the improved local lemmas in the later sections of the proof, record every dependence of density, progression length, dimension, and correlation on earlier parameters, and obtain a complete explicit iteration theorem.

Proposition V.12.1 specifies the final bookkeeping for one common type of increment. The hard work is to prove the hypotheses with the newly propagated functions. An initial useful deliverable would be a symbolic parameter dependency graph accompanied by Lean theorems for each monotone substitution. This would reveal which improvements survive and which are absorbed by later losses, even if the final asymptotic class stays unchanged.

Research question V.14.8 (Quantitative rounding and efficient derandomization). Optimize the group-size requirement for the sharpness constructions at a given error tolerance, including the exact-density version of Corollary III.2.12. Can one derandomize the rounding with an efficient finite algorithm?

The existence problem at a compatible prescribed rational density is already solved by the slow perturbation and small alteration in Corollary III.2.12. Its size requirement is deliberately generous. Sampling directly with a fixed total, a sharper norm estimate for the alteration, or a deterministic conditional-expectation argument might improve it. Independence no longer applies unchanged after conditioning on cardinality, so that approach needs its own estimate.

Research question V.14.9 (A reusable formal library of centered configuration counts). Build a finite-abelian-group library whose basic objects are systems of integer linear forms, centered multilinear averages, character constraint sets, and explicit reindexing bijections. Recover the cube and progression results as specializations.

The key interface should retain integer torsion data. A proof over $\mathbb{Q}$ alone does not control evaluation on all finite groups. At the same time, the library should make cancellation by a single free variable easy, so simple cases do not require a full Smith normal form calculation. This is an achievable software and formalization project with reusable mathematical outputs, independent of completing the entire Gowers proof.

V.15 Reproducibility and verification

Source 03, delivered Appendix A.

The accompanying archive contains the complete LaTeX source and compiled PDF, the figure in vector PDF form, scripts for the figure and finite checks, and the recorded check output. No original third-party paper is repackaged. The repository revision and source paths are recorded so that the comparisons can be repeated against the same input.

V.15.1 What the scripts check

`verify_cubes.py` uses only the Python standard library and exact rational arithmetic. It enumerates every four-vertex support of the Boolean cube for dimensions two through five. In dimension five this means all $\binom{32}{4} = 35960$ supports. It compares a classification from minors with a separate classification from XOR and coordinate patterns. The counts are exactly those in Lemma III.2.5. It also computes the coefficients of the full cube polynomial directly, without using the proposed quartic formula, on cyclic groups of orders two through six in dimensions two and three. The low-degree cancellations, torsion correction, and parity polynomial all pass.

`verify_estimates.py` checks the sharp selection extremizers with exact rational weights, evaluates the cubic certificate including its boundary and tangent cases, constructs proper phase partitions at several finite scales, and enumerates all subsets of $\mathbb{Z}/5\mathbb{Z}$ and $\mathbb{Z}/7\mathbb{Z}$ for the progression bounds at lengths three and four. The norm powers in the latter check are rational; taking roots and comparing the bounds uses floating-point arithmetic with an explicit tolerance. The phase oscillation checks are also numerical. These are clearly distinguished from the exact cube and scalar checks.

Independent written reviews checked the asymmetric selection extremizers, finite progression geometry, Cauchy–Schwarz induction, quartic support classification and Bernoulli rounding, cubic minorant, and four-label counting injection. These reviews and finite computations are useful error checks; the proofs in the article, rather than the tested finite range, justify the theorems for arbitrary parameters.

V.15.2 Building the document

From the extracted directory, run

```
python3 verify_cubes.py
python3 verify_estimates.py
pdflatex -interaction=nonstopmode -halt-on-error gowers_local_refinements.tex
pdflatex -interaction=nonstopmode -halt-on-error gowers_local_refinements.tex
```

The figure is already included, so compiling the document does not require Python or a plotting library. To regenerate the figure, run `make_figure.py` in an environment with NumPy and Matplotlib. The LaTeX source uses standard packages including AMS mathematics, Palatino fonts, `microtype`, `booktabs`, and `hyperref`.

V.16 Quantitative comparisons and what they enable

Source 04, delivered Section 9.

The preceding results address different sources of loss. Their benefits should therefore be read at the interface where each estimate is used. The following table records those interfaces without combining unrelated parameters into a single claimed global improvement.

The field-size requirement in the restriction theorem has a particularly simple interpretation:

$$|B| \geq \Lambda_d(\alpha\eta)^{-2d}.$$

Interface	Coarse source estimate	Strengthened estimate here
Graph moment, input $\gamma B ^3$	$\gamma^7\beta^{21}p^{15}$	$\gamma^7\beta^{15}p^{15}$; arbitrary moment order and equality cases
Almost-homomorphism restriction, order eight	$(\gamma^7\beta^6\eta/4)^{2^{19}} B ^{15}$	$2^{-95}\gamma^{112}\eta^{15} B ^{15}$, under the explicit size hypothesis in Corollary II.9.5
Three-cube upper bound, $u = \ f\ _{U^3}$	$(\delta + u)^8$	Excess begins with $6\sqrt{2}\delta^4u^4$, with explicit higher terms
k -term progression counting	Error bounded by $2^k u$	Error bounded by $u \sum_{j=0}^{k-3} \delta^j$
Section 17 local coefficient	$2^{-2(k+1)^3}$	$2^{-k^2 - k \log_2 k - O(k)}$, with an explicit constant

Table 11: Comparisons with the specified local estimates in [1, 23]. The variables and hypotheses are those of the corresponding sections. The improved restriction theorem includes a large-field condition, and the phase theorem includes a characteristic condition.

The ambient density does not otherwise enter this sufficient condition. Two elementary decisions are responsible. A rank-two exceptional relation leaves at most $|B|^{2d-2}$ tuples, because its free coordinates can be chosen from B . These contributions are then summed with their nonnegative Fourier weights, whose total is at most one. Thus the number of possible coefficient vectors does not enter the error bound. Repeated entries are handled separately. These refinements are useful even before any attempt to optimize the kernel.

For an illustrative comparison, take $\gamma = \eta = \beta = 1/2$. The retained fraction of $|B|^{15}$ in the source display is $2^{-8388608}$, whereas the convenient bound here is 2^{-222} . The corresponding sufficient field-size condition here is $p \geq 2^{231}$. These are deliberately transparent evaluations of the proved formulas, not practical finite-size estimates. They show the change in parameter dependence and also the cost of the conservative exceptional-tuple estimate.

The cube results answer two precise local extremal questions. The best universal leading coefficient in the real centered, odd-order setting is $6\sqrt{2}$. In the setting without 2- or 3-torsion, the first possible correction after that leading term has optimal exponent $16/3$. What remains undetermined is the exact coefficient of this correction and the existence of a universal asymptotic expansion for the extremal envelope. These are local conclusions with explicit models, independent of any claim about the best quantitative bound for arithmetic-progression-free sets.

[write] **6 October 2026; sources 11–14.** The exact coefficient is now determined: $(9/4)^{1/3}$, by source 14 (Remark W4).

V.17 Further research questions and a proposed program

Source 04, delivered Section 10.

The questions below are proposed research targets arising from the calculations. The list distinguishes exact optimization problems from larger integration projects. It is not a claim that each formulation is already recognized as an open problem in the literature; a more extensive priority review should accompany work toward publication.




Figure 2: Proved bounds and an exact extremizing model at $\delta = 1/4$. Left: the excess above δ^8 in the source triangle-inequality bound and the two refined bounds of Theorems III.3.6 and III.3.8. The refined curves nearly coincide at small u because they have the same sharp leading term. Right: after subtracting that leading term, the explicit two-frequency family of Proposition III.3.9, divided by $\delta^{8/3}u^{16/3}$, tends to $(9/4)^{1/3}$. The figure evaluates formulas; it is not a statistical fit.

V.17.1 Extremal cube counts

Q1. Does the secondary extremal coefficient exist, and what is it? Fix $0 < \delta < 1$, and let $M_\delta(u)$ be the supremum of $Q_3(W)$ over prime cyclic groups of order greater than 16 and weights $0 \leq W \leq 1$ satisfying $\mathbb{E}W = \delta$ and $\|W - \delta\|_{U^3} = u$. Theorems III.3.8 and Proposition III.3.9 imply

$$(9/4)^{1/3} \leq \liminf_{u \downarrow 0} \frac{M_\delta(u) - \delta^8 - 6\sqrt{2}\delta^4 u^4}{\delta^{8/3}u^{16/3}} \leq \limsup_{u \downarrow 0} \frac{M_\delta(u) - \delta^8 - 6\sqrt{2}\delta^4 u^4}{\delta^{8/3}u^{16/3}} \leq 6. \quad (\text{Q1})$$

Determine whether the limit exists. Is the two-frequency value $(9/4)^{1/3}$ optimal? A useful first step is a sharper estimate of the quintic correlation in terms of the exact loss in the norm-gap inequality, instead of the separate bounds used above.

[write] **6 October 2026; sources 11–14.** *Settled by source 14*, by a route close to the one suggested here (it sharpens the loss in the norm gap by eight further Fourier slots and balances it against the quintic correlation): the limit exists and equals $(9/4)^{1/3}$, and the two-frequency value is optimal (Theorem III.34.1; in this normalization, with prime cyclic groups of order above 16 and norm exactly u , Remark W4). Item III.3.

Q2. What is the correct stability statement for cube extremizers? Corollary III.3.5 controls the fourth-power Fourier mass outside one conjugate pair, equivalently a U^2 remainder. This does not alone control an L^2 remainder: many small Fourier coefficients can have little fourth-power mass but substantial second-power mass. Determine the strongest stability conclusion for bounded weights that nearly attain the leading cube bound, and identify any additional hypothesis needed for a stronger metric. For near optimal secondary behavior, test whether concentration on frequencies $\{\pm\xi, \pm2\xi\}$ follows.

[write] **6 October 2026; sources 11–14.** *Answered for the secondary scale by source 14*: near-optimal secondary behaviour forces concentration on $\{\pm\xi, \pm2\xi\}$ in U^2 , with amplitude and phase alignment (Theorem III.34.3), and the indicator examples show that no L^2 conclusion follows. Item III.4.

Q3. What replaces the sharp norm gap in higher dimensions? For $d \geq 4$, determine the best constant in the comparison of $\|f\|_{U^d}^{2^d}$ and $\|f\|_{U^2}^{2^d}$ for real centered functions, specifying the allowed group orders. Combine it with the exact coefficient P_d of Proposition III.3.2. The Boolean support matroid gives a finite organizing device for the next correlations, but their values and sharp norm bounds require additional harmonic analysis. Separate large-characteristic behavior from the extra identities produced by small torsion.

[write] **6 October 2026.** For $d = 3$ the comparison is Theorem III.3.4; for $d \geq 4$ this is source 03's Research question V.14.1 (which asks for the constant over odd-order groups). Merged in Section V.2, item III.1.

[write] **6 October 2026; sources 28–33.** *Advanced for $d \geq 4$, not settled* (sources 28 and 29, Section III.66). Source 29 proves $\|f\|_{U^d}^{2^d} \geq 2^{-m} \binom{2m}{m} \|f\|_{U^2}^{2^d}$, $m = 2^{d-2}$, hence $c_4^{\text{odd}} \leq (8/35)^{1/4}$ and $c_d^{\text{odd}} \leq B_d \downarrow 1/2$, and $c_d^{\text{odd}} \geq 1/3$ for every d (Corollary III.69.1, Proposition III.69.2); source 28 proves $c_4^{\text{odd}} \geq 0.5587294$ and that a real part of a character is not extremal for $d \geq 4$ (Corollary III.70.8, Theorem III.70.9); the intake's certificate gives $c_4^{\text{odd}} \geq 0.5587725$. So $0.5587725 \leq c_4^{\text{odd}} \leq (8/35)^{1/4} = 0.6914416$ among sources 01–33; later placed sources record better upper bounds (item III.1).

V.17.2 Restriction and moment estimates

Q4. Can the size threshold be substantially lowered? Theorem II.9.1 requires $|B| \geq \Lambda_d(\alpha\eta)^{-2d}$. The proof already controls exceptional relations in expectation using their Fourier weights. Can their contribution be made to reflect the additive structure of B , rather than the uniform rank-two bound? Can a kernel with a more favorable coefficient distribution reduce the required size without weakening the retained tuple count? Any proposed improvement must handle repeated entries as well as extra relations among distinct entries.

Q5. What is the true optimal retention law? The bound $\kappa_d \alpha^{2d} \eta^{2d-1} |B|^{2d-1}$ has the natural parameter powers for the single-kernel expectation budget in Proposition II.9.6. That proposition does not establish optimality among all selection procedures. Construct upper-bound examples, or find a different selection method with better powers. Also ask for useful simultaneous guarantees on $|B'|$, its ordinary additive energies, and the density of respected relations.

Q6. How far does restriction extend beyond prime cyclic groups? The interpolation theorem already applies to arbitrary finite abelian groups. The restriction theorem uses a field to turn two independent linear equations into a uniform bound on the number of solutions. Develop versions for $\mathbb{Z}/N\mathbb{Z}$, finite-field vector spaces, and maps into other groups. The appropriate hypotheses should quantify kernels of the relevant coefficient maps, rather than silently treat nonzero determinants as invertible. A weighted formulation would also allow nonuniform input measures.

Q7. Can the cross-section profile be propagated through the proof? For arrangements the gain is governed by $\sigma = \mathbb{E}_r \beta_r^3$, not just by the global density. Track this profile through subsequent restrictions and changes of variables. Determine which operations preserve useful control of σ , and whether a collection of higher cross-section moments gives a stronger interface than a single density parameter. Equality models in Section II.10 provide tests against overly optimistic exponents.

V.17.3 Patterns, counting, and integration

- Q8. Determine the asymptotics of the fractional pattern counts.** The proved estimates imply

$$\frac{1}{2} \leq \liminf_{k \rightarrow \infty} \frac{\log_2 F_k}{k^2} \leq \limsup_{k \rightarrow \infty} \frac{\log_2 F_k}{k^2} \leq 1, \quad (\text{Q2})$$

and the same bounds for A_k . Determine the leading constant, or improve the gap in these bounds. Relate the particular affine arrangements here to established results on threshold patterns and subset-sum arrangements before assigning novelty. Exact certified counts for $k = 4$ are a concrete computational next step; exploiting symmetries is likely to matter more than merely using a faster numerical solver.

[write] 6 October 2026. *Answered by source 01* (cross-source finding of the intake, checked at the write). The F_k here is source 01's C_k and the A_k here is source 01's A_k : the definitions are identical. Source 01's Theorem IV.2.3 gives $k^2 - k \log_2 k - O(k) \leq \log_2 C_k \leq k^2 + (\log_2 e - 1)k$. Its Theorem IV.3.1 gives $C_k \leq A_k$ and $\log_2 A_k \leq k^2 + O(k)$. Hence both limits in (Q2) exist and equal 1. The lower half uses one external input, the classical threshold-function count (Zuev; Kahn–Komlós–Szemerédi, as recorded by Baldi–Vershynin [8]), which source 01 cites and does not reprove. Source 01's Proposition IV.5.1 certifies $F_4 = C_4 = 154$, the “concrete computational next step” asked for. A_4 is not certified (an uncertified recount gives 770), and the secondary term is source 01's first question; both remain open (Section V.2, items IV.2 and IV.1).

[write] 6 October 2026; sources 06–10. Source 06 answers the same question independently: its Theorem IV.7.6 gives $\log_2 C_k = (1 + o(1))k^2$ and $\log_2 D_k = (1 + o(1))k^2$ (the same external threshold count), uniformly modulo every $M \geq 2$, and its Section IV.9 certifies $C_4 = 154$ a second time.

- Q9. Which global and local profiles are actually realizable?** The phase-extraction proof uses the uniform bound 2^k for the profiles at a fixed slope and an arrangement bound for all slopes. The slopes there come from a short discrete interval. Determine the joint distribution of profiles in that arithmetic setting, including small moduli M , and test whether an average multiplicity can replace the worst-case local multiplicity. The exact modular formula of Theorem IV.6.3 isolates one source of collisions.
- Q10. Can progression counting use more of the mixed information?** The dynamic program in Section III.4 finds the best bound within a specified telescoping family. It is not an optimality theorem for all counting arguments. Investigate whether mixed uniformity norms, additional cancellations, or information about particular pairs of factors can lower that bound. Any extension to nonadmissible slopes should display the relevant subgroup averages explicitly; the $\mathbb{Z}/2\mathbb{Z}$ example rules out simply dropping the invertibility assumptions.
- Q11. What survives complete propagation through the density increment?** Create a dependency graph for the quantitative parameters of the original proof. Substitute the improved moment, restriction, and profile estimates, and recompute every subsequent threshold with the same hypotheses. The output should be an explicit final function of k and density, with each loss traceable to a proved lemma. Compare the resulting bound with both the original argument and later work [2, 3]; an improvement in the historical proof need not improve the stronger modern theorem.
- Q12. Can the refinements be verified by a proof assistant?** Formalize the lowest-risk interfaces first: exact moment interpolation, the density-weighted telescope, and the finite support enumeration. Then formalize the selection theorem, where the main obligations are

finite Fourier orthogonality, independent sampling, and the exceptional rank count. A fully checked dependency chain would clarify which quantitative gains are compatible with the later induction, and would separate a correct theorem statement from a completed formal proof. Appendix V.19 gives a proposed implementation order.

Questions 1, 4, and 8 offer relatively contained starting points: a sharp analytic constant, a finite selection threshold, and a concrete arrangement census. Questions 7 and 11 are integration projects, for which local success must be accompanied by careful propagation of hypotheses. The distinction helps keep a proposed research program testable.

V.18 Closing perspective

Source 04, delivered Section 11.

The strongest conclusions of this investigation are explicit local theorems with complete proofs. Exact moment accounting feeds a polynomial restriction theorem; centered Fourier structure yields sharp cube coefficients and a fractional correction exponent; and simultaneous floor-pattern counting replaces a cubic exponential local loss by a quadratic one. The sharpness examples identify which aspects have already reached their natural limit and which constants remain to be determined.

The next substantial step is to turn these interfaces into a fully checked quantitative dependency chain. That would establish how much of the local improvement survives later stages of Gowers’s argument, while the extremal and pattern questions above offer independent research directions with precise targets.

V.19 A map to the Lean development

Source 04, delivered Appendix A.

This appendix identifies interfaces in the inspected snapshot [23]. It records a plan for additional work; it is not a report that new Lean declarations have been proved. A declaration defining a proposition, a theorem with a proof term, and a theorem successfully checked by the Lean kernel are different pieces of evidence.

One signature issue deserves explicit review. The inspected definitions of `lemma_9_3` and `corollary_9_4` in `Sections08_09.lean` quantify over nonzero moduli without an explicit primality typeclass. The selection theorem in this article uses primality essentially. It must therefore be formalized as a separate prime-specialized statement, or accompanied by a justified revision of the intended source interface. It is not automatically a proof of the broader existing proposition. The repository already has proofs of those statements through another route; the present claim is a stronger prime-field interface. By contrast, the exact moment inequality itself does hold over arbitrary finite abelian groups.

[write] 6 October 2026. Checked at HEAD: `lemma_9_3` and `corollary_9_4` still quantify over all $N \geq N_0$ without primality and have exact companions, as stated. The table omits one interface that was already formal at source 04’s pin: Lemma III.4.5 is `hasNatAP_of_short_modular_sequence` (note there). The other proposed strengthenings are not formalized.

For a practical implementation order, first isolate a normalized and an unnormalized Fourier convention with explicit conversion lemmas. Next prove the exact moment and telescoping statements,

Inspected interface	Proposed strengthening and principal proof obligations
<code>Proofs09Moments.lean;</code> <code>lemma_9_2_holds</code>	Keep the exact second moment before applying the higher-moment inequality. State the arbitrary-order graph version, with $ B $ in the denominator. The proof already computes an exact cardinality before weakening it.
<code>Proofs03CubeUpper.lean;</code> <code>lemma_3_10_holds</code>	Add centered cancellation and the real odd-order norm gap as separate lemmas. Prove the twelve quartic and eight quintic support classifications with explicit finite certificates or a transparent finite case split.
<code>Proofs03ProgressionAverage.lean;</code> <code>theorem_3_2_holds</code>	Generalize the required changes of variables to admissible slopes. Make multiplication automorphisms explicit, and reuse the Gowers–Cauchy–Schwarz infrastructure.
<code>Proofs03ProgressionCount.lean;</code> <code>corollary_3_3_holds</code>	Replace the full subset expansion by a finite telescoping identity. Retain density products and cancel the exact two-factor average.
<code>Proofs09Restriction.lean;</code> <code>lemma_9_3_holds</code> and <code>corollary_9_4_holds</code>	Introduce a prime-specialized version of the strengthened restriction theorem. Formalize the rank-two completion bound, the nonnegative Fourier weights, and the random-selection expectation.
<code>Sections17_18.lean</code>	Add definitions of fractional patterns and normalized profiles, followed by the hyperplane bound and the local phase-extraction theorem with explicit characteristic assumptions.

Table 12: Suggested interfaces for formalizing the proved refinements.

which need little specialized finite enumeration. Then introduce a library of finite Boolean support certificates for the cube identities. Finally add the probabilistic selection and floor-pattern arguments. Each step should compile against the pinned toolchain of the repository, and the final verification report should name the checked declarations and the commit. None of those future verification steps is represented as complete here.

V.20 Reproducibility and supplementary checks

Source 04, delivered Appendix B.

The source archive contains the complete L^AT_EX manuscript, its compiled PDF, the figure source and PDF, a provenance manifest, and three independent verification programs. The programs are deliberately supplementary to the proofs. In particular, testing finitely many functions cannot prove the norm-gap theorem, and checking small groups cannot prove the restriction theorem for all sufficiently large primes.

V.20.1 Energy, cube, and coefficient checks

`verify_refinements.py` uses the Python standard library. The recorded run checked all 8 470 nonempty graphs of partial maps $\mathbb{Z}/n\mathbb{Z} \rightarrow \mathbb{Z}/n\mathbb{Z}$ for $n = 2, 3, 4, 5$, giving 59 290 exact integer instances of the energy inequality for moment orders $2, \dots, 8$. It also enumerated the three-cube support ranks over the rationals and over $\mathbb{F}_2, \mathbb{F}_3, \mathbb{F}_5, \mathbb{F}_7$, and counted parallelograms through dimension seven. These computations exhibit the different support count in characteristic two.

The kernel checks use exact rational arithmetic for the selected coefficient ratios and the central-

binomial inequality, and exact integer arithmetic for the convenient powers of two in Corollary II.9.5. The complete two-frequency polynomial in Proposition III.3.9 is independently recovered with rational coefficients.

Floating-point checks are identified separately in the results file. They test the Fourier identity and norm gap on 40 real centered functions on small odd cyclic groups, the cosine examples, and 36 density-weighted progression identities, including an admissible example modulo 25. The largest recorded Fourier-identity discrepancy is less than 1.1×10^{-16} , and the progression identity discrepancy is less than 3.4×10^{-17} . The asserted numerical tolerance is 2×10^{-11} . These are consistency checks, not rational certificates for arbitrary functions.

The supplied `verify_selection.py` checks the selection expectation bounds directly by exact rational Fourier expansion in 10 examples over $\mathbb{F}_3, \mathbb{F}_5, \mathbb{F}_7$, covering 40 584 ordered additive tuples. It treats each distinct selected point only once and includes 504 distinct unrespected six-tuples. These checks use the degree-one kernel and supplement the proof of the general weighted selection lemma; the large-degree constants are checked separately.

V.20.2 Exact floor-pattern certificates

`pattern_census.py` verifies the supplied `pattern_census.json` using only rational arithmetic from the Python standard library. It obtains

$$(F_1, F_2, F_3) = (1, 2, 10), \quad (A_1, A_2, A_3) = (2, 6, 38).$$

Every branch in the hyperplane enumeration has one of two certificates: a rational point satisfying all relevant strict inequalities, or nonnegative rational weights whose combination proves that those strict inequalities are infeasible. The verifier checks both sides of every split, so it checks completeness as well as feasibility. There are 450 certified branches in the supplied record. Generating the certificates uses a numerical linear-programming solver to propose candidates, but accepting them and verifying the saved census does not use numerical tolerances.

V.20.3 Build and verification commands

From the extracted archive directory, the following commands rebuild the article and check its supplementary records:

```
latexmk -pdf -interaction=nonstopmode -halt-on-error \
  gowers_quantitative_refinements.tex
python3 -S verify_refinements.py --output verification_results.json
python3 -S verify_selection.py --report selection_verification.json
python3 -S pattern_census.py --verify pattern_census.json
```

The prebuilt vector figure is included, so rebuilding the article does not require Python plotting libraries. To regenerate the figure, run `python3 make_figures.py`; that optional step requires NumPy and Matplotlib. Certificate generation additionally requires SciPy. The source manifest records the inspected repository commit and the relevant Git object identifiers. The archive does not redistribute the original paper or the full repository.

V.21 Relation to the Lean development

Source 05, delivered Section 7.

V.21.1 Inspected source and a conservative integration plan

The relevant files are `Proofs02Partition.lean`, `Proofs02PhasePartition.lean`, and `Proofs02DensityIncrement.lean`; the first and third were inspected in detail in the ranges relevant to these results. The existing density-increment file imports the phase-partition file. The repository uses `NatAP`, its carrier and properness predicate, and a finite partition relation. Its partition construction already groups residue classes into blocks with lengths at least L and less than $2L$, matching the integer mechanism used here [25].

The efficient formalization route is to keep that construction and expose a parameterized theorem, rather than duplicate it for each numerical constant. The mathematical contract should assume $2L(L-1)M \leq rs$ in the modular version and conclude a partition with cell lengths in $[L, 2L-1]$ and image diameter at most s . The proof then invokes the existing small-image-step lemma with $Q = \lfloor r/L \rfloor$.

The sharp capped result becomes a separate integer calculation with $U = \lfloor \sqrt{rs/M} \rfloor$ and $K = \lceil U/2 \rceil$. The optimality example can be formalized independently using the cardinality of a disjoint union of two-element cells. It does not require complex analysis.

For the density theorem, the main additional interfaces are the exact identity $\sum |\mathbf{1}_A - \delta| = 2\delta(1-\delta)N$, a midpoint or disk-center phase estimate, and the capacity-sensitive positive-mass lemma. The existing file already proves the zero-sum identity, density bounds, positive/negative mass equality, and nonzero Fourier-coefficient bounds by both the set and its complement [26].

V.21.2 Recommended separation of proof layers

There are four natural layers. The first is integer combinatorics: proper progressions, residue classes, lengths, and disjointness. The second is phase approximation on $\mathbb{R}/\mathbb{Z}$ or $\mathbb{ZMod} N$. The third is finite weighted real and complex inequalities. The fourth specializes to Fourier coefficients and density increments. Keeping these layers separate permits the same weighted lemmas to be used with later polynomial-phase partitions.

For a minimum-risk first contribution, formalize the sharp $1/3$ modular theorem and the strengthened Corollary 2.4 before introducing the spectral-energy material. The simultaneous box lemma is another independent milestone. The algebraic obstruction is mathematically self-contained, but its Lean implementation would additionally need algebraic-number and resultant or norm infrastructure; it need not block the upper partition bounds.

V.21.3 Verification status

No new Lean file in this package is claimed to compile, and no new theorem has been kernel-checked here. The package supplies ordinary mathematical proofs, an exact-arithmetic implementation of the finite partition constructions, and a formalization roadmap. We did not perform a whole-repository proof audit or infer the status of the complete Gowers development from the inspected modules.

The observed `main` reference during this work was

327773149bd272239f1f6c0015f76d5097512c88.

The accompanying source-audit file records the content hashes of the compared files. The partition module was also fetched explicitly at that commit. The other inspected contents are identified by their returned blob hashes; the moving branch is not treated as a permanent identifier.

[write] **6 October 2026.** Checked at HEAD: the cited blobs are correct and unchanged (291b149 for `Proofs02Partition.lean`, 85b5c3a for `Proofs02DensityIncrement.lean`, 41085a0 for the transcription). One defect of the shipped audit file `05-phase-partitions-SOURCE_AUDIT.md`: it says the commit-pinned read of `Proofs02Partition.lean` “included lines 580–780”, but that file has 702 lines; the blob and the construction it describes (blocks of lengths at least L and less than $2L$) are correct. The audit file is shipped unchanged (README).

V.22 Reproducible computation and the limits of testing

Source 05, delivered Section 8.

The companion implementation `code/phase_partitions.py` uses Python integers and rational fractions for all partition decisions. It implements the exact scale (I.16.2), the capped $1/3$ theorem, the simultaneous ceiling condition (I.18.2), and the tolerance allocation (I.19.10). An independent checker verifies complete coverage, no repeated points, proper arithmetic progressions, one common difference, the claimed cell lengths, the number-of-cells bound, and exact circular arc widths.

The completed exact tests comprise 13,325 one-phase cases, 11,440 capped cases, 320 simultaneous cases, and 200 tolerance-allocation cases. The one-phase tests vary the interval length from 1 to 41, all slopes with denominators from 2 to 11, and five rational widths. Of the simultaneous cases, 166 have minimum length greater than one. The seeded random tests use dimensions from one to four. These checks test rounding and implementation; they are not substitutes for the proofs in Sections I.16 and I.18.

Fourier diagnostics use 80 decimal digits. There are 60 correlation-and-popularity cases and 96 spectral-energy cases. Twenty-four spectral cases have a non-singleton guaranteed minimum length, including six genuinely multifrequency cases at larger sizes. The larger examples are modular pullbacks of intervals, with their Fourier coefficients computed by a geometric-series formula independently of the partition construction. All tested inequalities passed with positive observed margins. These are numerical diagnostics, not exact certificates for transcendental inequalities.

The verification files distinguish exact structural tests from high-precision diagnostics and explicitly record that new Lean verification was not performed. To reproduce the results, run the commands in the package README. No source paper, repository file, or font file is redistributed.

V.23 Further research questions

Source 05, delivered Section 9.

The following questions arise naturally from the proved results. Unless explicitly stated otherwise, they are research directions formulated here; we do not assert that they are all previously recognized open problems.

1. Exact finite partition profiles. The universal constant $1/3$ with the fixed upper cap is settled above, but the exact best lower length for each triple (r, s, M) and each slope can be much

larger. Determine an efficient characterization of the feasible pairs of minimum and maximum lengths, taking account of residue-class sizes rather than only a worst-case blocking lemma. The explicit obstruction at total size nine explains one universal constant, not the full profile.

2. Best constants in the simultaneous minimax theorem. The exponent and the symmetric dependence on ε are sharp, but the upper constant obtained from the elementary algebraic norm argument is very poor. Determine the optimal leading constants, first for one or two phases, and the extent to which requiring a common difference across the whole partition changes them. One must specify the order of the limits in N and ε .

3. Anisotropic lower bounds. The lower construction depends on $\prod_i \varepsilon_i$. Is there a matching minimax upper bound in all highly anisotropic regimes, up to dimension-dependent constants and an additive one? A coordinate with a very loose tolerance can cease to be a genuine constraint, so a naive product formula may require a regime decomposition. Weighted Diophantine approximation is the natural language, but no such theorem is assumed here.

4. Optimal density-sensitive correlation transfer. The disk radius and the capacity inequality are separately sharp, but their simultaneous extremizers need not arise from an indicator and a linear phase. Determine the exact feasible region of block positive mass, Fourier correlation, density, and phase width. Even a finite-dimensional extremal formulation could improve the trade-off (I.17.9) and identify which constant losses are genuine.

5. Adaptive selection of frequencies and tolerances. The continuous allocation problem is solved by Proposition I.19.3, but selection of the frequency subset and optimization of the exact ceiling formula are not. Develop an algorithm that balances captured energy, coefficient imbalance, dimension, and the final integer progression length. Its correctness should be stated relative to a certified Fourier input rather than floating-point estimates alone.

6. Higher-degree and nilsequence interfaces. Apply the residual-variance inequality to valid higher-degree localization theorems. The challenge is to track density-sensitive oscillation budgets together with geometric complexity and exceptional cells. A useful target is an explicit theorem where all losses are parameterized, so that an improvement of a partition result propagates mechanically through the density argument. This would still need a separate analysis before it implied any global Szemerédi improvement.

7. Formalizing sharpness as well as existence. A Lean formalization should include the parity obstruction proving optimality of $1/3$, not just the existence theorem. This makes the quantitative API resistant to later accidental strengthening. The next independent target is the exact residual inequality (I.19.4); it isolates a reusable finite-probability fact without depending on the full higher-order proof.

8. Interval-localized spectral extraction. On a proper interval, restricted characters have a nontrivial Gram matrix. Corollary I.19.4 provides an explicit scalar-norm correction, but a large value of κ can erase the gain from several correlations. Establish useful certified conditioning bounds for specific frequency configurations, and optimize the linear combination using more of the Gram

matrix than its largest eigenvalue. This would turn the abstract extension into an effective local Fourier theorem.

V.24 Conclusion

Source 05, delivered Section 10.

The finite localization step admits stronger statements than the corrected repository currently exports. The lower constant $1/3$ in the capped partition theorem is exactly optimal; the new Corollary 2.4 retains the same upper cell size with a better lower size and correlation retention; and the density increment uses the true mass $2\delta(1 - \delta)$ instead of a unit bound. Simultaneous localization has a clean, sharp exponent, and a residual-variance inequality turns an appropriate Fourier-energy hypothesis into a multiplicative increment with a density-sensitive tolerance budget.

These are local quantitative results with complete proofs and reproducible computational checks. Their place in the larger program is to provide stronger interfaces and clearer bottlenecks, not to bypass the higher-order structural work needed for a new global theorem.

V.25 An explicit potential for density-dependent iteration

Source 06, delivered Section 6.

The iteration in the proof of Gowers' Theorem 18.2 estimates every density increment by the one available at the starting density [1]. Tracking the increasing density instead gives a sharper bound. We first prove a general deterministic lemma; it separates this bookkeeping improvement from the local structural theorem that supplies the increment.

Lemma V.25.1 (Discrete potential). *Let $0 < \delta_0 \leq 1$ and $t \geq 1$. Let F be positive, nondecreasing, and continuously differentiable on $[\delta_0, 1]$. Let G be nonnegative and nonincreasing there. Suppose*

$$\delta_0 < \delta_1 < \cdots < \delta_{t-1} \leq 1, \quad \delta_{i+1} \geq \delta_i + F(\delta_i).$$

Then

$$\sum_{i=0}^{t-1} G(\delta_i) \leq G(\delta_0) + \int_{\delta_0}^1 G(u) \left(\frac{1}{F(u)} + \frac{F'(u)}{F(u)} \right) du. \quad (\text{V.25.1})$$

Proof. For $x < y$ with $y - x \geq F(x)$, put $r = F(y)/F(x) \geq 1$. Monotonicity gives

$$\begin{aligned} \int_x^y \left(\frac{1}{F(u)} + \frac{F'(u)}{F(u)} \right) du &\geq \frac{y-x}{F(y)} + \log r \\ &\geq \frac{1}{r} + \log r \geq 1. \end{aligned}$$

The last inequality follows because $r^{-1} + \log r$ is nondecreasing for $r \geq 1$ and equals one at $r = 1$. Since $G(u) \geq G(y)$ on $[x, y]$, the same integral weighted by G is at least $G(y)$. Apply this observation to the interval preceding each summand except the first, and add. $\square$

Taking $G = 1$ gives a useful count alone:

$$t \leq 1 + \int_{\delta_0}^1 \frac{du}{F(u)} + \log \frac{F(1)}{F(\delta_0)}. \quad (\text{V.25.2})$$

For $F(u) = cu^p$ with $p > 1$, this becomes

$$t \leq 1 + \frac{\delta_0^{1-p} - 1}{c(p-1)} + p \log \frac{1}{\delta_0}. \quad (\text{V.25.3})$$

The logarithmic correction makes the integral estimate valid for the discrete process without assuming infinitesimal steps.

V.25.1 A power-law density increment and length loss

Theorem V.25.2 (Explicit iteration budget). *Suppose a hereditary configuration problem has the following valid local alternative. At any density $0 < \delta \leq 1$ and length $N \geq N_* > 1$, either the desired configuration exists, or restriction to a subprogression gives density and length satisfying*

$$\delta' \geq \delta + c\delta^p, \quad N' \geq N^{a_0\delta^q}, \quad (\text{V.25.4})$$

where $c > 0$, $0 < a_0 \leq 1$, $p \geq 1$, and $q \geq 0$ are fixed. Assume the same alternative applies after each restriction, and fix an initial density $0 < \delta_0 \leq 1$. Set

$$A = \log(1/a_0), \quad L = \log(1/\delta_0).$$

For $p > 1$, put $r = p - 1$ and define

$$\begin{aligned} B(\delta_0) = A + qL + \frac{A\delta_0^{-r} - 1}{c} \frac{1}{r} \\ + \frac{q\delta_0^{-r}(rL - 1) + 1}{c} \frac{1}{r^2} + pAL + \frac{pq}{2}L^2. \end{aligned} \quad (\text{V.25.5})$$

For $p = 1$, define instead

$$B(\delta_0) = A + qL + \left(\frac{1}{c} + 1\right) \left(AL + \frac{q}{2}L^2\right). \quad (\text{V.25.6})$$

Then a sufficient initial length is

$$\log \log N \geq \log \log N_* + B(\delta_0). \quad (\text{V.25.7})$$

All logarithms in these formulas are natural.

Proof. Writing $H = \log \log N$, the length conclusion in (V.25.4) loses at most

$$G(\delta) = A + q \log(1/\delta)$$

from H at one step. Apply Lemma V.25.1 with $F(u) = cu^p$. Its integral is

$$\int_{\delta_0}^1 (A + q \log(1/u)) \left(\frac{1}{cu^p} + \frac{p}{u} \right) du. \quad (\text{V.25.8})$$

For $p > 1$, the elementary identities

$$\int_{\delta_0}^1 u^{-p} du = \frac{\delta_0^{1-p} - 1}{p-1}$$

and

$$\int_{\delta_0}^1 u^{-p} \log(1/u) du = \frac{\delta_0^{1-p}((p-1)L-1) + 1}{(p-1)^2}$$

evaluate (V.25.8) and give (V.25.5). Direct integration for $p = 1$ gives (V.25.6).

If no desired configuration exists, apply the local alternative repeatedly. After every finite number of restrictions, (V.25.1) and (V.25.7) leave $\log \log N_i \geq \log \log N_*$, so the local alternative remains applicable. Every step increases density by at least $c\delta_0^p > 0$. Such a sequence cannot stay at most one indefinitely, which is a contradiction. $\square$

A nonincreasing density-dependent threshold $N_*(\delta)$ causes no change: use $N_*(\delta_0)$ as the common threshold. A fixed length prefactor is also explicit. If the local length conclusion is

$$N' \geq C^{-1} N^{a_0 \delta^q}, \quad C \geq 1,$$

use the stronger working threshold

$$N_{\text{work}} = \max \left\{ N_*, \exp \left(\frac{2 \log C}{a_0 \delta_0^q} \right) \right\}. \quad (\text{V.25.9})$$

At every retained density, this implies

$$\log N' \geq a_0 \delta^q \log N - \log C \geq \frac{a_0 \delta^q}{2} \log N.$$

Thus Theorem V.25.2 applies with a_0 replaced by $a_0/2$ and N_* replaced by N_{work} .

V.25.2 Optimal leading order and the scope of the improvement

For fixed a_0, c, p, q with $p > 1$, formula (V.25.5) has the expansion

$$B(\delta) = \frac{\delta^{1-p}}{c} \left(\frac{A + q \log(1/\delta)}{p-1} - \frac{q}{(p-1)^2} \right) + O((\log(1/\delta))^2 + 1). \quad (\text{V.25.10})$$

Its leading term is optimal for the abstract local rules (V.25.4). Indeed, take the slowest allowed density sequence

$$\delta_{i+1} = \delta_i + c\delta_i^p$$

and charge exactly $G(\delta_i)$ at each step until a fixed positive target δ_* is reached, choosing $\delta_* + c\delta_*^p < 1$ so that the last step stays below 1. Since G/F is nonincreasing,

$$\int_{\delta_i}^{\delta_{i+1}} \frac{G(u)}{F(u)} du \leq \frac{G(\delta_i)}{F(\delta_i)} (\delta_{i+1} - \delta_i) = G(\delta_i).$$

Summing gives a lower bound by the main integral, up to a constant depending on the fixed target. Lemma V.25.1 exceeds that integral by only $O((\log(1/\delta))^2 + 1)$. Consequently, the leading coefficient in (V.25.10) cannot be improved using only the local recurrence and its charged length loss. Additional structural information would be necessary.

For comparison with the scale in Gowers' final iteration, suppose a valid local theorem supplies

$$\delta' \geq \delta + (\delta/2)^p, \quad N' \geq N^{(\delta/2)^p}. \quad (\text{V.25.11})$$

Then $c = a_0 = 2^{-p}$ and $q = p$. For fixed $p > 1$ and $\delta \rightarrow 0$, the leading double-logarithmic cost is

$$\frac{2^p p}{p-1} \delta^{1-p} \log(1/\delta).$$

Freezing density throughout gives a cost asymptotic to

$$2^p p \delta^{-p} \log(1/\delta).$$

Their ratio is asymptotic to $\delta/(p-1)$: the potential saves one power of the starting density in the double-logarithmic threshold. This changes neither the tower height nor the local structural input. In particular, the repository snapshot explicitly records that its edited proof of Theorem 18.1 no longer establishes the displayed tower exponent after its revision of Corollary 16.11 [41]. Specializing (V.25.11) to that particular exponent is therefore conditional on a valid local increment theorem; Theorem V.25.2 itself has the complete proof given above.

The two refinements in these sections are complementary. The extraction theorem preserves the information already present in a partition, including the occupied negative mass and any discarded exceptional set. The iteration theorem preserves the subsequent increase in density. Both identify exactly which additional information would be required for a larger improvement.

[write] **6 October 2026; sources 06–10.** Iteration bookkeeping is also in source 03’s Proposition V.12.1 and source 10’s Propositions V.42.1 and V.42.2; the leading coefficients of the step counts are compared in the note after source 03’s proposition (the coefficient here is the smallest). The statement about Theorem 18.1 is still true at HEAD: `theorem_18_1`, `theorem_18_2` and `corollary_16_11` are open in the ledger.

V.26 Compatibility with the Lean development

Source 06, delivered Section 8.

The new mathematical statements in this report have not been added to the repository or compiled in Lean. The purpose of this section is to give an implementation map grounded in inspected definitions and proof endpoints, so that a subsequent formalization has a precise starting point. All names below belong to the namespace `LeanProofs.GowersSzemerédi` in the pinned snapshot [41].

V.26.1 Normalization and statement boundaries

The repository’s `gowersNorm d f` is unnormalized:

$$\|f\|_{U^d(G)} = \frac{\text{gowersNorm } d \ f}{N^{(d+1)/2^d}}.$$

Its predicate `UniformOfDegree f alpha d` corresponds to

$$\|f\|_{U^{d+1}}^{2^{d+1}} \leq \alpha.$$

Both the normalization and the index shift must be preserved when translating the quadratic counting criterion. The development’s `progressionAverage` is an unnormalized double sum, so the normalized discrepancy formulas here acquire a factor N^2 .

The existing endpoint `theorem_3_2_holds` lives in `Proofs03ProgressionAverage.lean`. The exact source companion `corollary_3_3_holds` is in `Proofs03ProgressionCount.lean`;

it currently retains the 2^k factor. The source’s monotonicity and progression-existence companions are in `Proofs03Basic.lean` and `Proofs03ProgressionExistence.lean`. The public `gowers_cauchy_schwarz` in `Proofs03GowersCS.lean` supplies the mixed-cube inequality used in the localized phase proof.

V.26.2 A practical sequence of formalization steps

1. **Introduce a generic weighted extraction theorem.** Formalize Theorem I.21.1 using finite positive weights and cell averages, independently of cyclic groups. Derive `lemma_5_15_holds` and its stronger replacements as special cases. The proof needs finite-sum identities, non-negativity, and pigeonhole arguments. The rational extremizers provide explicit sharpness witnesses.
2. **Replace the subset expansion in progression counting.** First add the indexed telescoping theorem without changing the existing uniformity assumptions. Then prove the two-target Cauchy–Schwarz induction and the quadratic bound. Keeping the first-order theorem separate is useful because it needs uniformity only for the later inputs, whereas the convenient homogeneous quadratic statement assumes a common higher norm bound.
3. **Expose finite carry families.**
The source defines `floorDotPattern` and `floorAffinePattern`. The two exact companions are `corollary_17_5_holds` and `corollary_17_6_holds`. Their proofs appear in `Proofs17FinitePatterns.lean`. That file splits the intercept into integer and fractional parts but does not use the exact split of every slope developed here. New finite-image definitions for $C(V)$ and $D(V)$ would make Theorems IV.7.3 and IV.7.4 natural cardinality statements. The one-sided perturbation lemma should be formalized before invoking the region count.
4. **Parameterize localized phase removal by a pattern bound.** The proof in `Proofs17LocalizedPhaseRemoval.lean` already isolates the raw inequality with a factor L^2 and a numerical lemma requiring $L^2c < 1/4$. Expose a version whose input is any valid pattern bound $L \leq K$; instantiate it with $c = 1/(8K^2)$. This permits a local replacement without rewriting the Fourier and cube algebra.
5. **Keep the Section 18 dependency conditional until its inputs are proved.** The actual theorem `section18_cyclic_density_increment_holds_of_theorem_18_1` has an explicit Theorem 18.1 hypothesis. Writing its exponent as ε , it obtains a density gain $\varepsilon/4$ and a progression size at least $(\varepsilon/4)N^\varepsilon$. The size prefactor must be retained or absorbed by a proved threshold estimate. The potential in Section V.25 is a separate, reusable way to perform the remaining iteration once a stable local theorem with explicit hypotheses is available.

The interval and affine transfer modules already document an additional quantitative issue: embedding a progression of length L into a cyclic group of order M changes ambient density by the factor L/M . A proof that repeatedly uses the new modulus cannot keep the old density without accounting for this factor. Our modular existence criterion does not remove that obligation.

V.26.3 What a completed formalization would certify

The existing audit distinguishes ordinary theorem companions from conditional implications and checks their axiom dependencies. A formalization of the proposed refinements should retain that distinction. There are three different goals: proving each auxiliary inequality, proving that its hypotheses hold at a particular use site, and proving that all constants survive the complete iteration. The first two can be made substantially more transparent by the exact interfaces above; the third remains a separate quantitative theorem.

[write] **6 October 2026; sources 06–10.** *Checked at HEAD (764740f07):* every declaration named above exists, with the roles described (`theorem_3_2_holds`, `corollary_3_3_holds`, `corollary_3_6_holds`, `gowers_cauchy_schwarz`, `corollary_17_5_holds`, `corollary_17_6_holds`, `lemma_5_15_holds`, `section_18_cyclic_density_increment_holds_of_theorem_18_1` in `Proofs18QuantitativeSzemerédi.lean`). None of this source’s statements is formalized.

V.27 Further research questions and a concrete program

Source 06, delivered Section 9.

The questions below arise directly from the proved results. Unless a separate literature status is given, “question” means that the issue is not settled in this report; it does not assert that no answer has ever appeared elsewhere. This qualification is particularly relevant to small chamber counts and to inequalities obtainable by standard Cauchy–Schwarz arguments.

- 1. Exact carry enumeration and a combinatorial interpretation.** Find a structural formula or recurrence for C_k and D_k , beginning with the certified values $C_4 = 154$ and $D_3 = 38$. The arrangement has normals indexed by nonempty subsets, integer thresholds, and a distinguished unit cube. Its symmetries, intersection lattice, and periodic extension are natural tools. An OEIS investigation should identify the underlying objects and definitions, not merely match a short numerical prefix. The exact certificates here provide a small benchmark for a symmetry-aware enumerator.
- 2. The next term in carry entropy.** The leading coefficient in $\log_2 C_k \sim k^2$ is now determined by the threshold reduction and arrangement bound. Determine the order and coefficient of the next term, or prove a substantially smaller upper/lower gap. The current elementary upper bound is $k^2 + O(k)$; a sharper use of the threshold embedding or of the geometry internal to the unit cube may distinguish C_k from the unrestricted threshold count.
- 3. Quantify information lost modulo a fixed modulus.** Although $C_{k,M}$ has the same leading entropy as C_k for every $M \geq 2$, this leaves open a large range of possible ratios $C_{k,M}/C_k$. Determine that ratio, or its logarithmic scale, for $M = 2$ and $M = 3$. The exact factorization (IV.7.6) means that any such result immediately improves modular pattern bounds without changing the integer-part argument.
- 4. An exact affine orbit formula.** The nonaffine decomposition is injective because singleton carries vanish. Affine carries do not have that property. Classify precisely when two fractional affine carry maps differ by an integer linear map modulo M , and use this classification to determine $L_k(r, M)$ in the saturated range $2r \geq M$. This is a finite equivalence problem that can be explored with the certified carry families before seeking an asymptotic theorem.

5. **Avoid the square of the pattern count analytically.** The proof of Theorem IV.8.1 uses one Cauchy–Schwarz step over classes and one choice of a class, giving L^2 . Is there useful orthogonality, an energy-weighted decomposition, or additional structure among the realized classes that reduces this loss? The entropy lower bound rules out subquadratic enumeration of all classes, but it does not rule out an analytical argument that pays only for the classes carrying significant energy.
6. **Optimal density dependence in higher progression counts.** For $k = 3$, the variance factor in the quadratic discrepancy bound is sharp, as singletons show. For $k \geq 4$, determine the best coefficient depending on k and δ that can replace $\sqrt{\delta}B_k(\delta)$. One can first ask the narrower question within the family of repeated telescoping and mixed-norm Cauchy–Schwarz proofs. This would separate improvements of the method from a sharp extremal theorem about all sets.
7. **Torsion-sensitive forms of the two-target inequality.** When multiplication by a slope difference has a kernel, the variable changes in the proof cease to be measure-preserving bijections. Determine the best replacement constants in terms of those kernels and their interactions. Three-term progressions are the first test case; longer systems may require data beyond individual kernel cardinalities. This would extend the present theorem to moduli excluded by $\gcd(N, (k-1)!) = 1$.
8. **Use the mass of all dense cells.** The sharp extraction theorem quantifies the total mass above each threshold, not just the size of one selected cell. Can a branching density-increment argument exploit this mass to improve a structural or counting bound? The extremizers show that an improvement from the same scalar data alone is impossible; one must retain geometric information about the collection of high-density cells.
9. **Optimal finite corrections to the iteration potential.** For the power-law recurrence, the leading integral cost is already optimal. Determine a minimal discrete correction, or an exact dynamic-programming potential, with a useful closed form for finite starting density. A companion problem is to incorporate a density-dependent length prefactor and a nonmonotone admissibility threshold without replacing them everywhere by their worst initial values.
10. **Close the quantitative dependency chain in the inspected formalization.** The highest-impact repository task is to complete the outstanding structural estimates around Sections 13 and 16, establish the needed form of Theorem 18.1 with explicit thresholds, and only then combine extraction, transfer, and iteration. The local improvements here should be retained as parameters rather than immediately weakened to one coarse tower expression. A second route is to plan a formal treatment of modern inverse-theorem inputs, making clear which substantial theorems are being imported.
11. **Other integer-valued configuration families.** Theorem IV.7.3 already handles finite non-negative integer coefficient families V . Study which families have carry entropy comparable to the rank times the logarithm of their size, and which have additional cancellations. For signed families, first resolve the lower-dimensional floor-boundary obstruction identified after that theorem. Such a generalization could support localized arguments for configurations other than a Boolean cube.

For immediate implementation, the weighted extraction theorem and the indexed telescoping estimate are the smallest independent formalization tasks. The exact carry decomposition is the next natural target; it has an elementary proof, useful finite certificates, and a direct downstream

consumer in phase removal. A global improvement requires the additional structural inputs identified above, with their actual density and scale dependence preserved.

[write] **6 October 2026; sources 06–10.** These eleven questions are merged into Section V.2 with credits (items IV.2, IV.1, IV.3, IV.4, III.8, III.9, I.11, V.3, V.1 and IV.6). Question 7 is partly advanced by source 10’s Theorem III.28.2, which handles kernels for one distinguished function.

V.28 Verification and the Lean integration boundary

Source 07, delivered Section 9.

V.28.1 Exact finite checks actually performed

The supplementary `checks.py` uses only Python’s standard library. All decisions use integers or `fractions.Fraction`; the sole decimal parameter `ratio` is for display. The recorded run produced `verification.json`, with all checks passing.

Check	Completed instances
All maps $\phi, \psi : \mathbb{Z}/4\mathbb{Z} \rightarrow \mathbb{Z}/3\mathbb{Z}$, all order-two source relations	6,561 function pairs; 419,904 relations
Equal-total word pairs in 220 unequal-fiber or punctured examples, using both the original and a second starting law	24,200 relation–law pairs
Primal profile versus exact one-variable dual	16,472 equalities
Nested-to-average envelope, including empty targets	9,059 applicable cases
Collision-to-mode inequality	1,936 cases
Sharp cyclic winding family, $2 \leq k \leq 20$, $2 \leq m \leq 5$	76 examples
General allowance at $\eta = (40k)^{-5}$, $2 \leq k \leq 1000$	999 rational endpoint checks

The randomized finite suite has the fixed seed 20261006 and includes both unstructured values and nearly affine examples. It recorded 6,058 strict relation certificates, so its tests are not confined to vacuous large-error bounds. The exhaustive singleton suite recorded 5,880 strict certificates and 157,464 violated relations, each with budget at least one. These are finite sanity checks on normalization, strictness, and source-law handling. They do not establish any universal statement by themselves; the general proofs are Sections II.14–II.19.

The inequalities for the order-five, order-nine, and relaxed order-two endpoints were also evaluated exactly. For example, the first budget at $\eta = 360^{-5}$, $\theta = 1/36$, is

$$18 \left(\frac{71}{1296} + 360^{-5} \right) = \frac{331257600001}{335923200000} < 1.$$

The original hypothesis is strict, but checking this closed endpoint gives a slightly stronger numerical verification of the corollary.

[write] **6 October 2026; sources 06–10.** Rechecked at the second write in exact arithmetic: $q(1/36) = 71/1296$, $18(71/1296 + 360^{-5}) = 331257600001/335923200000 < 1$, and $360^{-5} = (40 \cdot 9)^{-5}$, which is why Corollary II.17.6 follows from Theorem II.17.5 at $k = 9$.

V.28.2 What was inspected in the repository

The baseline is commit `ae0e895748209f0dd079c6b5ca6df6c6fc065a81`. The relevant files and roles are:

`Section10.lean`

Its local-model definition records the two nested agreement conditions. Its shift-regular definition includes fiber saturation, symmetry, factor-two fiber comparability, and overlap. The catalogue entry `lemma_10_9` concludes `FreimanHom 2 B' psi`.

`Proofs10Shift.lean`

The file contains `lemma_10_9_holds`. Its proof includes explicit handling of nonempty terminal fibers; this note does not attribute a vacuity bug to that proof. The proposed refinement replaces its two-step exceptional-set organization with averaged path marginals.

`Audit.lean`

Its comments distinguish an axiom-boundary audit from discharging the hypotheses of conditional results and from matching the numbered catalogue to the paper. This distinction is retained here.

The statements and proof source were read, but no Lean build or audit was run for this article. In particular, the existence of the old companion proof is not evidence that the new order-five or order-nine theorem has been formalized. The references [42, 43, 44] and the machine-readable source manifest identify the comparison precisely.

[write] **6 October 2026; sources 06–10.** At HEAD (764740f07) `Section10.lean`, `Proofs10Shift.lean` and `Audit.lean` are unchanged since the pin `ae0e89574`; `lemma_10_9` is at line 335 and `lemma_10_9_holds` at line 717. The ledger has changed since the pin (4 statements closed: `lemma_16_1`, `lemma_16_5`, `lemma_16_6` and `lemma_16_9`), none of them in Section 10.

V.28.3 A concrete formalization route

A practical integration should first prove a finite counting version of the four-edge theorem and only then introduce arbitrary paths and optimized profiles. A proposed module decomposition is:

1. **Directional kernels.** Define internal target fibers, boundary counts, and nonnegative rational edge defects. Prove the identities (II.13.3)–(II.13.4), with positivity hypotheses attached only to denominators actually used.
2. **Transport bounds.** Prove the simple domination lemma $\sum \mu_s f \leq c(s) \sum \mu f$ for $0 \leq f$, and derive the factor-two and additive-invariance instances. The full profile optimization can be postponed; it is unnecessary for the five/nine corollaries.
3. **A shared-endpoint finite sample.** Represent paths by finite lists or `Fin`-indexed vertices, with two explicit endpoint variables and independent internal variables. Prove a weighted union bound before translating it into existence of a successful sample.
4. **The local interface theorem.** Prove Lemma II.17.1, then the factor-two order- k criterion and the order-five consequence. Keep these separate from the stronger theorem using `DomainInvariant` and a fiber lower bound.

5. **Exact constants and audits.** Derive the real-power comparison once, discharge the rational polynomial estimates, and run the namespace axiom audit on the resulting new theorem declarations. Do not relabel an unproved proposition-valued catalogue entry as a completed theorem.

Two technical design choices deserve explicit tests. First, when a path returns to its starting index, its terminal *sample* must not be identified with its initial sample; otherwise the intended pair marginals can change. Second, restricting to previously successful equations is not the same as restricting to surviving indices. Only the latter restriction is used in the marginal bound. The file `FORMALIZATION.md` in the package records these obligations separately from completed mathematics.

V.29 Further research questions and priorities

Source 07, delivered Section 10.

The following are proposed problems arising from this framework. Except for the sharp balanced threshold already settled in Theorem II.16.1, we make no claim that these are previously catalogued open problems or that the listed approaches resolve them.

Question V.29.1 (Optimal nonuniform threshold). *For fixed $k \geq 2$ and $\kappa > 1$, what is the largest universal δ -threshold forcing order- k relations when fiber sizes vary by at most κ ? How does a separate boundary bound improve it?*

The bound $1/\{2[1 + (k - 1)\kappa]\}$ is proved, but only its balanced specialization is shown sharp. A promising finite search would optimize fiber sizes and actual value assignments, rather than arbitrary independent error profiles. The latter relaxation may exaggerate what a single ϕ can realize simultaneously.

Question V.29.2 (The strongest order at the original interface). *What is the maximal guaranteed Freiman order under the precise assumptions of the recorded Lemma 10.9, and under the fuller ambient-invariance package?*

Five and nine are lower bounds on the guaranteed order, not maximality results. The exact nested envelope, the actual value of ρ , optimized word ordering, and nonuniform directional defects can all improve the sufficient certificate. Counterexamples must also satisfy the nested local model, not merely a scalar defect bound.

Question V.29.3 (Stability of winding obstructions). *If constant fibers and a violated order- k relation have total budget close to one, must their error sets resemble a small collection of cuts or carry boundaries after an appropriate structural reduction?*

The cyclic construction has a concrete wraparound boundary. Near equality in the union bound suggests that edge-failure events should be almost disjoint in a shared-endpoint experiment, but translating that observation into structural information about ϕ requires additional work.

Question V.29.4 (Efficient certificates for many relations). *Can the root-law linear programs for a large family of relations be compressed into a small reusable set of certificates, perhaps by exploiting translation symmetries or a presentation of the generated group?*

Theorem II.18.3 gives a correct finite optimization for one relation. It does not give an efficient algorithm when the number of relations is exponential in the number of directions. A useful implementation would report rational primal witnesses whenever it certifies a relation.

Question V.29.5 (Joint optimization of extraction and transport). *Can one construct the regular component and its source law together, so that the extracted domain is large and its relevant transport profiles are small, instead of first paying for coarse fiber comparability?*

Changing the source law affects the hypothesis under which local errors are controlled. The alternative-root result avoids silently changing that hypothesis, but it does not solve the joint extraction problem. A successful method would need to account explicitly for lost cardinality and for the relation between the original energy weights and the chosen law.

Question V.29.6 (Higher-dimensional coherence). *Can shared-vertex sampling on cubes or more general finite diagrams replace repeated nested exceptional-set arguments in bilinear and multilinear exactification steps?*

The two-path proof suggests a diagrammatic error budget, but additional compatibility equations can correlate the required marginals. The right problem is to specify a joint sampling law first and prove every face or edge marginal under that law, rather than assume independent tests can be glued automatically.

Question V.29.7 (Improving the earlier extraction exponent). *Which quantitative losses before the local model are necessary, and can the appearance of $\eta^{1/5}$ itself be improved while retaining a sufficiently large regular component and direction set?*

The present method removes the later square root; it does not improve the earlier fifth-root construction. This is a more consequential target than optimizing the last few units in a rounded denominator such as 75.

Question V.29.8 (An end-to-end quantitative payoff). *After substituting the transport-sensitive exactification theorem, which later bottleneck becomes dominant, and does any improvement remain in a fully tracked final bound?*

Answering this requires a complete dependency ledger through the later Bohr-set, multilinear, and density-increment steps. The initial deliverable should be a mechanically checkable parameter recurrence, followed by a comparison with the unchanged proof, before any improved global exponent is announced.

Recommended order of work. First formalize the factor-two order-five theorem and its explicit boundary handling. Next build an exact rational profile-certificate tool and search for sharp nonuniform examples. Only then undertake the full parameter propagation and higher-dimensional extensions. These tasks separate near-term checkable improvements from the substantially larger research program.

[write] 6 October 2026; sources 06–10. The eight questions are items II.9–II.15 and V.1 of Section V.2.

V.30 Conclusion

Source 07, delivered Section 11.

The local algebraic obstruction is a closed diagram with inconsistent edge labels. Once a joint law with shared endpoints is fixed, a violated relation must consume at least one unit of total failure probability. Transport profiles express exactly how an error known under the original law can contribute along a translated path. This yields linear, rather than square-root, error propagation and preserves useful information about higher-order relations.

The concrete consequences are a sharp balanced threshold, an order-five strengthening at the recorded Lemma 10.9 interface, an order-nine conclusion when ambient invariance is retained, and larger explicit local order-two allowances. The article proves those claims and supplies finite exact-arithmetic checks. Their significance is local and structural; an improved global Szemerédi bound would require the additional end-to-end analysis identified above.

V.31 Reproduction commands

Source 07, delivered Appendix B.

From the unpacked directory, the mathematical checks and PDF can be reproduced with

```
python3 checks.py --output verification.json
latexmk -pdf -interaction=nonstopmode -halt-on-error transport_rigidity.tex
```

A Makefile provides `make check`, `make pdf`, and `make clean`. The document uses an embedded bibliography and does not need BibTeX. Recompilation requires the LaTeX packages listed in its preamble. No network access is required for either command once those packages are installed.

V.32 Integration with the supplied Lean development

Source 08, delivered Section 11.

V.32.1 The normalization boundary

The inspected definitions use $\mathbb{Z}/N\mathbb{Z}$, complex-valued functions, the unnormalized Fourier transform, and the unnormalized cube norm [45]. Translating the mathematical notation gives

$$\begin{aligned}\text{cubeCount}(d, A) &= N^{d+1} \mathcal{C}_d(\mathbf{1}_A), \\ \text{gowersNorm}(d, f) &= N^{(d+1)/2^d} \|f\|_{U^d}, \\ \text{UniformSetOfDegree}(A, \alpha, d-1) &\iff \|\mathbf{1}_A - \delta\|_{U^d}^{2^d} \leq \alpha \quad (d \geq 1).\end{aligned}$$

The last equivalence uses the successor-dimension identity for the iterated difference sum. It is already the normalization step used by the existing proof of Lemma 3.10.

For $d \geq 2$, the proposed replacement conclusion is therefore

$$\text{cubeCount}(d, A) \leq N^{d+1} \left[\delta^{2^d} + \sum_{r=4}^{2^d} \binom{2^d}{r} \delta^{2^d-r} \alpha^{r/2^d} \right]. \quad (\text{V.32.1})$$

The dimension restriction avoids both the U^1 seminorm issue and the truncated natural subtraction in the expression $d-1$ at $d=0$. Dimensions zero and one already have elementary exact counts and should remain separate cases rather than being forced through this theorem.

V.32.2 A proposed proof decomposition

No new Lean file is included as if it were verified. The following are proposed lemma boundaries, not claims that these identifiers already exist. They are intended to let the existing analytic development carry most of the proof.

Suggested lemma or module	Mathematical obligation
<code>cubeForm_subset_expansion</code>	Expand the finite product into a sum over subsets of Boolean vertices.
<code>cubeProjection_three_uniform</code>	Build the affine equivalence implementing the unit-minor argument for three distinct vertices.
<code>centeredCubeTerm_eq_zero</code>	Factor averages for subsets of cardinality one, two, or three.
<code>centeredCubeTerm_norm_le</code>	Pad omitted vertices with 1 and apply the existing cube Cauchy–Schwarz inequality.
<code>cubeCount_fourthOrder_le</code>	Combine cancellation, the cardinality strata, and the normalization factor N^{d+1} .
<code>cubeQuartic_torsion_formula</code>	Add the Fourier classification of four-vertex subsets and the two closed counting formulas.

A useful implementation order is to prove the first four statements for the current cyclic setting before generalizing the entire API to finite abelian groups. The elementary fourth-order improvement does not require Fourier analysis, classification of finite abelian groups, or any torsion assumptions. The exact quartic formula is naturally a second, separate layer.

For the indicator theorem, the function `balanced A` is real-valued inside `C`. Recording that property explicitly lets parity conjugations simplify before the subset expansion. The existing cube argument uses a minus sign, while the counting convention uses a plus sign; negating all increments gives the required measure-preserving equivalence.

The proof of the analytic upper bound can be isolated from real-power arithmetic by first proving a statement in a nonnegative parameter t with a hypothesis on the normalized balanced norm. Only at the final step need one substitute $t = \alpha^{1/2^d}$. That separation reduces the chance that a bookkeeping error in roots or normalization obscures the combinatorics.

V.32.3 What was inspected and what was not verified

The two main source files used for this integration analysis were `Definitions.lean` and `Proofs03CubeUpper.lean`. Their inspected Git blob identifiers begin with `97113b7afa69` and `c017515bdd43`, respectively; complete identifiers and source URLs are in the accompanying provenance file. The latter file proves `lemma_3_10_holds` by applying Minkowski [46]. The analysis is not a repository-wide axiom audit.

No Lean executable was available in the working environment, and no Lean build or kernel verification of the new results was performed. The repository was not modified. The deliverable contains mathematical proofs, exact finite checks, and this integration plan, rather than an untested formalization presented as complete.

[write] **6 October 2026; sources 06–10.** At HEAD (`764740f07`) `lemma_3_10_holds` is still proved in `Proofs03CubeUpper.lean` (line 82) by the Minkowski bound; the statements of this source are not formalized.

V.33 Exact checks and reproducibility

Source 08, delivered Section 12.

The accompanying file `verify_cube_refinements.py` uses only Python's standard library. All its asserted comparisons use integers or `fractions.Fraction`; it does not approximate trigonometric values or compare floating-point cube counts.

The following checks were run successfully in preparing this memorandum. First, the program enumerates every four-element subset in Boolean dimensions 2 through 6, verifying the R_d, T_d, F_d formulas on a total of 673,227 subsets. It also checks the unit-minor property on all 5,580 triples in dimensions 2 through 5.

Second, it classifies all three-cube five-element supports into 48 containing a parallelogram and eight exceptional supports, and it verifies the missing-pair multiplicities 12, 12, 4 for six-element supports.

Third, it computes the complete coefficients of $\mathcal{C}_3(\delta + \eta f)$ by direct averaging for explicit integer-valued, mean-zero functions on cyclic groups of orders 3 through 9. This verifies the vanishing of degrees one through three, the all-group quartic formula, and, in odd order, the full three-cube identity. Additional quartic checks use the groups $(\mathbb{Z}/2\mathbb{Z})^2$, $(\mathbb{Z}/2\mathbb{Z})^3$, $\mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/3\mathbb{Z}$, and $(\mathbb{Z}/3\mathbb{Z})^2$. These cover examples with several order-two characters as well as odd-order noncyclic groups.

Finally, two independent Fourier enumerations verify (III.10.3). One enumerates all eight frequencies; the other uses the four-parameter kernel formula (III.10.4). For a single cosine on $\mathbb{Z}/11\mathbb{Z}$, the same exact method gives

$$\mathcal{C}_3(\delta + \eta \cos(2\pi x/11)) = \delta^8 + \frac{3}{2}\delta^4\eta^4 + \frac{1}{32}\eta^8. \quad (\text{V.33.1})$$

The order-two quartic coefficient is checked algebraically in dimensions 2 through 7.

To rerun the tests from the extracted archive, use

```
python verify_cube_refinements.py --output verification_results.json
```

The default enumeration stops at dimension six. The optional argument `-max-d 5` shortens the quartic enumeration. Running the program rewrites only the chosen output JSON file; it makes no network requests and does not access the repository.

These checks are useful for catching coefficient, sign, and torsion mistakes. They are not proofs of statements quantified over all groups or all dimensions. Those statements are supported by the arguments in the preceding sections. The finite counterexample identity has its own explicit exhaustive certificate, whose relation to the Fourier equations is given above.

V.34 Further research questions

Source 08, delivered Section 13.

The following questions are proposed directions arising from the present results. Unless explicitly connected to a cited source, their appearance here does not assert that they are established open problems in the literature. The priority order favors tasks that can produce a concrete theorem or a checkable counterexample.

V.34.1 Optimal odd-order constants

Question V.34.1. *For $d \geq 3$, determine*

$$\kappa_d^{\text{odd}} = \sup_{\substack{G \text{ finite, odd order} \\ f: G \rightarrow \mathbb{R}, \mathbb{E}f=0, f \neq 0}} \frac{E_4(f)}{\|f\|_{U^d}^4}.$$

Is $\kappa_3^{\text{odd}} = 1/\sqrt{2}$?

[write] **6 October 2026; sources 28–33.** *Advanced for $d \geq 4$, not settled* (sources 28 and 29, Section III.66).

Source 29 proves $\|f\|_{U^d}^{2^d} \geq 2^{-m} \binom{2m}{m} \|f\|_{U^2}^{2^d}$, $m = 2^{d-2}$, hence $c_4^{\text{odd}} \leq (8/35)^{1/4}$ and $c_d^{\text{odd}} \leq B_d \downarrow 1/2$, and $c_d^{\text{odd}} \geq 1/3$ for every d (Corollary III.69.1, Proposition III.69.2); source 28 proves $c_4^{\text{odd}} \geq 0.5587294$ and that a real part of a character is not extremal for $d \geq 4$ (Corollary III.70.8, Theorem III.70.9); the intake's certificate gives $c_4^{\text{odd}} \geq 0.5587725$. So $0.5587725 \leq c_4^{\text{odd}} \leq (8/35)^{1/4} = 0.6914416$ among sources 01–33; later placed sources record better upper bounds (item III.1).

[write] **6 October 2026; sources 34–39.** *The upper end improved again* (sources 34 and 35, Sections III.71–III.72). Source 34 proves $\|f\|_{U^2}^4 \leq (2/3) \|f\|_{U^4}^4$, that is the coefficient $81/16$ in source 29's normalization (Theorem III.71.1); source 35, the same day, proves the coefficient 5 (Theorem III.72.1); for $5 \leq d \leq 8$ source 34's coupled bounds slightly improve source 29's (Remark III.71.10). So $0.5587725 \leq c_4^{\text{odd}} \leq 2/3$ among sources 01–39; source 45, placed and not yet written, records $(832/4795)^{1/4} = 0.6454070$ (item III.1).

[write] **7 October 2026; sources 40–46.** *The upper end improved again* (source 45, Section III.89): $\|f\|_{U^4}^{16} \geq (4795/832) \|f\|_{U^2}^{16}$ (Theorem III.90.1), so $0.5587725 \leq c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070$ among sources 01–46; source 49, placed by f8bc5e2ec and not yet written in, records $29/4$, that is $c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ (for its write) (item III.1).

[write] **7 October 2026; sources 47–53.** *The upper end improved again* (source 49, Section III.92): $\|f\|_{U^4}^{16} \geq (29/4) \|f\|_{U^2}^{16}$ (Theorem III.93.1), so $0.5587725 \leq c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ among sources 01–53, and $c_5^{\text{odd}} \leq 29/50$ (item III.1).

Norm monotonicity gives $\kappa_d^{\text{odd}} \leq 1$. The single-cosine calculation (V.33.1) gives $\kappa_3^{\text{odd}} \geq 1/\sqrt{2}$. The best odd-order leading coefficient in a fourth-order bound purely in t is $R_d \kappa_d^{\text{odd}}$, not automatically R_d . A useful first step is a constrained Fourier optimization on odd cyclic groups, followed by a proof of any proposed inequality. Small-group optimization alone would not settle the supremum over all groups.

[write] **6 October 2026; sources 06–10.** *Answered for $d = 3$.* This κ_d^{odd} is source 03's c_d^{odd} . Source 04's Theorem III.3.4 gives $\kappa_3^{\text{odd}} = 2^{-1/2}$, attained by real parts of characters (proof in the note after that theorem), so the answer to the last sentence of the question is yes. For $d \geq 4$ it is open (Section V.2, item III.1).

Question V.34.2. *What is the largest constant c such that*

$$\mathcal{D}_3(F) \geq c \delta^4 E_4(F - \delta)$$

for every odd-order group and every $0 \leq F \leq 1$ with $\delta = \mathbb{E}F > 0$?

The results here give $2 \leq c \leq 15337/1296 < 12$. The lower endpoint follows from Theorem III.12.1; the upper endpoint follows from the explicit weight in (III.10.5). Optimizing the two-frequency family gives a slightly stronger upper bound by minimizing $12 - 2s + 6s^2 + s^4$, but there is no reason to assume this family is globally extremal. A proof separating local and global optimizers would already clarify the geometry of the cube polynomial.

[write] **6 October 2026; sources 06–10.** The failure of the constant 12 is kept on record as Remark W1; the question is item III.10 of Section V.2, merged with source 09's question on a corrected lower bound.

V.34.2 Higher coefficients and arithmetic dependence

Question V.34.3. *Compute the complete fifth- and sixth-order coefficients for general d , including their dependence on the torsion of G .*

The quartic coefficient only sees order-two characters. For larger supports, integer minors can have additional prime factors, and Fourier relations can couple several multiples of a frequency. A concrete program is to classify support matrices up to cube symmetry, compute their integer relation lattices, and produce exact character-sum profiles. The desired output is a proved formula with orbit counts, not just a table of numerical values in sample groups.

Question V.34.4. *Can the survivor polynomial $\sum_r A_{d,r}(G)z^r$ be computed substantially faster than enumerating all subsets of V_d ?*

Proposition III.14.2 reduces much of the large-prime case to column dependence. Symmetry reduction, deletion–contraction ideas, and certified integer normal forms may make dimensions four and five tractable. Small-characteristic exceptions must be retained rather than inferred from the rational calculation. It would be valuable to output compact certificates that a proof assistant can check independently of the enumeration engine.

V.34.3 Stability and explicit constructions

Question V.34.5. *Describe near-extremizers for the quartic coefficient $R_d E_4(f) + T_d E_{\langle 2 \rangle}(f)$ relative to $\|f\|_{U^d}^4$.*

On groups with an order-two quotient, the exact extremal model supplied here is a real character. A stability theorem would quantify whether a nearly optimal ratio forces concentration on one character or on a small collection of characters. On odd groups, the correct model is not yet determined here. This question concerns a small-norm perturbation of a constant, rather than near-maximal U^d norm under an L^∞ bound; those are distinct extremal problems.

Question V.34.6. *Derandomize the indicator sharpness construction while keeping its fourth-order asymptotic coefficient.*

The rounding proof is existential and uses very large prime groups. A deterministic construction should guarantee rounding error $o(\eta^4)$ in U^d , not merely small discrepancy against individual Fourier characters. The collision estimate gives a clear quantitative target against which an explicit construction can be measured.

V.34.4 Localization, sparse weights, and formalization

Question V.34.7. *Which exact cancellations survive when cube averages are localized to intervals, progressions, or Bohr neighborhoods?*

Global three-vertex independence is the key input to the fourth-order bound. Restricting the parameter domain destroys that independence in general. A localized theorem would need either specially compatible averaging measures or explicit boundary and correlation errors. Establishing such a result is a more plausible route toward density-increment applications than directly substituting the present global estimate into an unrelated step.

Question V.34.8. *Can the inhomogeneous estimate be transferred to sparse majorized weights with an explicit error budget?*

For a relative counting theorem, the means and the majorant cannot be treated as interchangeable constants. One should first expand around the actual majorant model and identify which low-order correlations are controlled by its pseudorandomness assumptions. The finite character-kernel criterion may then organize the remaining balanced terms, but it does not supply the required majorant estimates by itself.

Question V.34.9. *Can the elementary fourth-order improvement and the quartic torsion formula be formalized as separate verified modules in the supplied development?*

The first task needs finite product expansion, affine equivalences, and the existing cube inequality. The second adds finite Fourier analysis and the classification of four vertices. A useful milestone is a theorem with the original unnormalized conclusion (V.32.1), together with tests that instantiate both odd moduli and even moduli. General finite abelian groups can be treated after the normalization and subset-expansion interfaces are stable.

V.34.5 From vanishing supports to sharper configuration bounds

Question V.34.10. *For a fixed non-cube system of linear forms, can one combine the exact full-support criterion with a true-complexity estimate to obtain the correct first nonzero power of the balanced uniformity parameter?*

This would separate two sources of quantitative loss: terms that should be zero before any norm inequality is used, and terms that are genuinely present but estimated inefficiently. The work of Gowers–Wolf and Manners [29, 14] supplies relevant norm-control theory, but the leading-power problem depends on the particular system and its arithmetic. It should not be assumed that arithmetic progressions inherit the same three cancellations as a full additive cube.

V.35 Conclusion

Source 08, delivered Section 14.

The main gain is a sharp change of scale in the error term of the selected cube-count lemma. Centering removes the first three perturbative orders, and the first surviving order is exactly controlled by parallelograms and order-two Fourier mass. The resulting exponent $4/2^d$ is optimal even for sets in odd prime cyclic groups.

The more detailed three-cube calculation reveals a second lesson: the next coefficient is signed. A positive fourth-order expansion does not imply a global lower bound with the same coefficient. The explicit two-frequency counterexample makes that obstruction quantitative.

These results leave the global Szemerédi problem unchanged, but provide a more precise local component, a finite-support classification framework, and concrete next problems with a path to computational and formal checks.

V.36 Build and source provenance

Source 08, delivered Appendix B.

The source document is self-contained LaTeX with an internal bibliography. From the extracted archive, compile it with

```
pdflatex -interaction=nonstopmode -halt-on-error fourth_order_cube_counts.tex
pdflatex -interaction=nonstopmode -halt-on-error fourth_order_cube_counts.tex
```

Two passes resolve references and the table of contents. The archive also includes `build.sh`, the exact-arithmetic verifier, its captured JSON output, a README, and `source_provenance.json`. It does not redistribute the original Gowers paper or any font files.

The original paper was checked at printed pages 483 and 485, corresponding to PDF pages 19 and 21 in the linked 124-page copy. The repository source inspection was read-only and took place on 6 October 2026. The provenance file identifies inspected blobs; it does not claim that the entire repository was frozen or audited at a single commit.

V.37 Reproducible exact checks

Source 09, delivered Section 11.

The supplied computations use exact integers and rational numbers. They are tests of the formulas and finite certificates, not a substitute for the all-dimension classification proofs.

V.37.1 Direct finite-group cube averages

The standard-library script `verification/verify_expansion.py` was run successfully on 24 specified examples. It covers cyclic groups of orders two through nine in dimension three, several cyclic groups in dimensions two and four, and product groups including $(\mathbb{Z}/2\mathbb{Z})^2$, $\mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/3\mathbb{Z}$, and $(\mathbb{Z}/3\mathbb{Z})^2$. For each case it constructs a deterministic integer-valued mean-zero function and enumerates every parameterized cube.

The coefficient of δ^{m-s} is computed directly as the s th elementary symmetric polynomial of the m evaluated function values, averaged over cube parameters. The script compares the first five coefficients with the formulas in Theorem III.21.1. This avoids both numerical roots of unity and Fourier rounding error. For the right-hand side it evaluates equivalent convolution formulas, including the quotient averages on $G/2G$ and $G/3G$ for the torsion statistics. The recorded file `results.json` contains every input and exact output.

The same script checks the Stirling formulas for $2 \leq d \leq 14$ and all $5^8 = 390625$ frequency assignments needed for the two-frequency polynomial in Appendix III.26.

V.37.2 Exhaustive small-cube lattice census

The optional SymPy script `verification/verify_lattices.py` was also run successfully with SymPy 1.14.0. It uses exact Smith normal forms for every four- and five-vertex subset in dimensions three and four. The five-vertex census is particularly informative:

Five-vertex type	$d = 3$	$d = 4$
Parallelogram plus independent fifth vertex	48	1200
Full-support real circuit (star)	8	160
Full rank, index one	0	2672
Full rank, index two	0	320
Full rank, index three	0	16
Total	56	4368

The index-two five-vertex configurations are a useful warning: their lattices are not primitive, but their moments still vanish because a character coordinate is forced to be trivial. Counting only nonprimitive supports would overcount the fifth-order contribution.

For four-vertex subsets the surviving counts are $(12, 2)$ in dimension three and $(100, 40)$ in dimension four, matching (P_d, T_d) . The same script checks the determinant, Smith invariants, and full modulo-five kernel certificate in Theorem III.23.1. Its complete output is `lattice_results.json`.

V.37.3 What was and was not verified

Both scripts completed with all assertions passing. The manuscript proofs handle arbitrary dimension and arbitrary finite abelian groups; the scripts do not exhaust those infinite families. No Lean compiler was available in the execution environment, no new theorem was checked by Lean, and no repository files were modified. The supplied formalization roadmap is therefore a proposal, not a verification certificate.

V.38 Integration with the supplied Lean development

Source 09, delivered Section 12.

V.38.1 Pinned source audit

The inspected repository checkpoint is commit

`7d99e5deeb209a4671060990cb1ea3d672d57064`.

The relevant directory is

`Combinatorics/Ramsey/Lean/GowersSzemerédi`.

The following are observations from source inspection, not claims that we rebuilt the repository:

Definitions.lean

Defines `fourier` as the unnormalized discrete Fourier transform, `balanced` as indicator minus density, `cubeForm` with parity conjugation and the argument $s - e \cdot x$, `gowersNorm` as an unnormalized cube-sum root, and `cubeCount` by parameterized additive cubes. The inspected blob is `97113b7afa6925a2dd4b76641eeaef09597ab6a`.

Sections01_03.lean

Contains the proposition-valued statement `lemma_3_10`. A proposition-valued catalogue entry alone is not a proof; here there is also a separate proof file.

Proofs03CubeUpper.lean

Proves `lemma_3_10_holds`. It imports `Proofs03Minkowski` and `Proofs03Cubes`, identifies indicator and constant norm powers, and applies `gowersNorm_add_le`. The inspected blob is `c017515bdd43b3c7dd8a5d0e90747a20403d4ddc`.

Proofs01_03.lean

Contains the finite Fourier identities, including Parseval and the fourth-moment/additive-quadruple identity. These are natural inputs for the exact quartic theorem.

V.38.2 Normalization and sign conventions

For $G = \mathbb{Z}/N\mathbb{Z}$, the dictionary is

Manuscript quantity	Repository quantity
$\widehat{f}(r)$	N^{-1} <code>fourier f r</code>
$\ f\ _{U^d}$	$N^{-(d+1)/2^d}$ <code>gowersNorm d f</code>
$\mathcal{C}_d(\mathbf{1}_A)$	$N^{-(d+1)}$ <code>cubeCount A</code>
$\ f_A\ _{U^d}^{2^d} \leq \alpha$	<code>UniformSetOfDegree A</code> α $(d - 1)$, for $d \geq 2$

The manuscript uses $x + v \cdot h$ whereas `cubeForm` uses $s - e \cdot x$. Replacing every side variable by its negative converts between them. Real-valued functions make parity conjugation harmless, but a formal proof must still discharge the corresponding real-cast and conjugation identities. The indicator count includes degenerate parameterized cubes; it is not a count of distinct vertex sets.

V.38.3 A staged implementation plan

The most economical first formal target is Theorem III.17.2, not the entire lattice classification. It needs only a subset expansion, three-vertex uniformity, and the existing Gowers–Cauchy–Schwarz infrastructure.

A proposed first module, `Proofs03BalancedExpansion.lean`, would prove that mixed cube sums supported on one, two, or three distinct Boolean vertices vanish for a real mean-zero function. Prove explicit coordinate reindexing equivalences rather than relying on field linear algebra over $\mathbb{Z}\text{Mod } N$. This avoids an unnecessary primality assumption. Then sum by subset cardinality and recover (III.17.2). A safe initial theorem should assume $d \geq 2$ and $\alpha \geq 0$ explicitly; this avoids accidental dimension-zero or real-power edge cases.

A second module, `Proofs03QuarticCircuits.lean`, would classify four-element Boolean subsets, count the two types, and prove their Fourier formulas. It can initially state a torsion-aware theorem for arbitrary N ; an odd-modulus corollary then removes the order-two term. Using an integer matrix or explicit coordinate-block normal form is essential. A theorem asserting independence over $\mathbb{Z}\text{Mod } N$ merely from rational rank would be false.

A third module, `Proofs03QuinticCircuits.lean`, would implement the rank-four and rank-five cases separately. The two especially important obligations are the primitive-minor argument for a

parallelogram plus a fifth vertex, and the index-two case with a forced trivial Fourier coordinate. The general character-kernel formula could be developed either first as a reusable library theorem or afterwards as a unification of the small cases.

A final `Proofs03CubeRefinement.lean` would combine these results with explicit binomial tails. The output theorem should coexist with `lemma_3_10_holds`, with the old theorem recovered as a corollary by adding back nonnegative omitted terms. There is no reason to delete a valid existing interface.

V.38.4 Recommended proof obligations and audit discipline

Useful public lemmas would include joint uniformity of at most three Boolean affine forms over an arbitrary finite abelian group; Fourier orthogonality for an integer matrix of forms; conversion of cube sums to normalized averages; counting disjoint nonempty coordinate blocks; and the binomial-tail inequality (III.16.5).

These should be genuine proved lemmas, not axioms or proposition placeholders. After implementation, build the exact pinned dependency set, audit the types of the new public theorems, inspect the axiom closure, and run the repository’s own status checks. A successful small experimental script is not a substitute for these steps, and the present package does not claim they have been carried out.

[write] **6 October 2026; sources 06–10.** At HEAD (764740f07) the inspected files still play the roles described; `lemma_3_10_holds` is at line 82 of `Proofs03CubeUpper.lean`. None of this source’s statements is formalized.

V.39 Further research questions and a prioritized program

Source 09, delivered Section 14.

The following are precise continuations of the proved results. They are proposed research targets, not assertions that each formulation is historically new or currently unsolved in the entire literature.

V.39.1 Complete the sixth-order coefficient

Research question V.39.1 (A dimension-uniform sextic expansion). Classify all six-vertex Boolean subsets that admit a full-support character tuple, and express the coefficient of δ^{m-6} in a finite basis of Fourier or higher-order statistics with closed formulas for their multiplicities as functions of d .

Theorem III.23.1 proves that an order-five term must occur. There will also be real relations and possibly several independent character parameters. A useful first computational stage is an orbit census by coordinate-pattern multisets, followed by integer-lattice certificates for each orbit. The real challenge is a finite normal-form classification that does not grow with the ambient dimension. The goal should be an exact sixth coefficient and a seventh-order remainder, not merely a longer numerical table.

V.39.2 Determine the prime-onset function

Research question V.39.2 (Support-efficient torsion). Determine $h(7)$ and $h(11)$, and improve the general bounds between (III.23.4) and $h(\ell) \leq \ell + 2$.

[write] **6 October 2026; sources 34–39.** *Answered by source 34:* $h(7) = 7$, $h(11) = h(13) = h(17) = 8$ (Theorem III.76.2), $h(p) = D_{\text{div}}(p) + 1$ (Theorem III.75.2) and $h(p) \sim 2 \log p / \log \log p$ (Theorem III.79.1); item III.16.

A determinant divisible by ℓ is necessary but not sufficient: the corresponding modulo- ℓ kernel must have full support, or balancing makes the term vanish. This extra support condition should be included from the start in a search. A practical certificate would contain the binary matrix, a nonzero determinant, its Smith invariants, and an explicit full-support kernel vector. The exact order-five example supplies a model for such a certificate.

V.39.3 Optimize constants at a fixed density and norm ratio

Research question V.39.3 (The constrained extremal envelope). For fixed d , density δ , and prescribed $(\varepsilon, \tau, \rho)$, determine or tightly bound

$$\sup\{\mathcal{C}_d(F) : 0 \leq F \leq 1, \mathbb{E}F = \delta, \|F - \delta\|_{U^d} = \varepsilon, \|F - \delta\|_{U^2} = \tau, \|\widehat{F - \delta}\|_{\infty} = \rho\}.$$

The coefficient A_d is optimal when all groups are allowed, but the unsigned fifth-order estimate need not be jointly sharp with it. Order-two characters maximize the universal fourth coefficient while forcing the quintic term to vanish. This incompatibility suggests that a coupled bound may improve (III.21.7). The odd-order problem is especially natural: determine the best leading coefficient relative to ε^4 , rather than relative to the exactly known Fourier energy τ^4 .

V.39.4 Stability of nearly extremal quartic corrections

Research question V.39.4 (Structure near the universal fourth-order maximum). If

$$P_d \mathbf{E}_4(f) + T_d \mathbf{T}_2(f) \geq (1 - \eta) A_d \|f\|_{U^d}^4,$$

what quantitative structural conclusion follows as $\eta \rightarrow 0$?

The order-two character example is an exact extremizer, but uniqueness and stability have not been proved here. One must understand both near-equality in norm monotonicity and concentration on the two-torsion spectrum. It would be useful to separate a statement valid for all real functions from stronger statements using $0 \leq \delta + f \leq 1$.

V.39.5 A useful corrected lower bound

Research question V.39.5 (Replace the false quartic lower bound). Find natural hypotheses under which M_5 is nonnegative, or bound its negative part more sharply in terms of the fourth-order statistics and density, so that a nontrivial corrected cube lower bound follows.

Nonnegative Fourier coefficients make all subset moments nonnegative, but this is a strong restriction. A less restrictive condition on the relative phases of a_χ and a_{χ^2} may be sufficient. The explicit expression (III.21.3) identifies the relevant phase interaction, and Proposition III.22.2 shows why an unconditional positivity argument cannot work.

V.39.6 Boundary-sensitive interval analogues

Research question V.39.6 (Cubes restricted to an interval). Develop a version of the expansion for cubes constrained to lie in $[N]$, with explicit boundary weights, and determine which of the first three cancellations survive exactly or approximately.

The current proof averages over a whole finite group. Conditioning on all vertices lying in an interval destroys the uniformity argument used in Lemma III.17.1. A direct transplantation would be incorrect. One possible route is to expand relative to a weighted reference measure adapted to the interval, rather than to a constant function, and explicitly retain the resulting lower-order boundary terms.

V.39.7 Algorithmic use of the measured residual

Research question V.39.7 (A residual-sensitive inverse step). Can the sixth-order certificate (III.21.11), together with an efficient estimate of the four Fourier statistics, improve a concrete inverse or counting algorithm in a useful parameter regime?

The question requires a complete cost and parameter analysis. A norm lower bound by itself does not deliver structure. A convincing result would specify the input access model, the cost of estimating $E_4, T_2, M_5, T_{3,5}$, the residual threshold, and the quantitative gain after an available inverse theorem is invoked.

V.39.8 Formalize primitive Boolean affine maps first

Research question V.39.8 (A reusable Lean library for arithmetic configurations). Build a library of exact reindexing equivalences and integer-lattice certificates for Boolean affine forms, then formalize Theorems III.17.2 and III.21.1 over arbitrary finite abelian groups.

The initial cancellation theorem is an accessible first milestone. The general character-kernel identity and a sound bridge from Smith normal form to finite-group character constraints would have much broader reuse. It is important to preserve non-prime groups in the interface: they expose the precise arithmetic phenomenon that the paper identifies.

V.39.9 Beyond diagonal cube counts

Research question V.39.9 (Colored cubes and other linear systems). For other finite systems of integer linear forms, identify the least support size at which a mean-zero expansion can contribute, and classify the first torsion-sensitive terms.

The proof mechanism is the same as Proposition III.18.1, but the first nonzero order need not be four. The heterogeneous theorem suggests applications to colored configuration counts and joint distribution tests. A useful result should state the primitive-lattice hypotheses explicitly, rather than replacing them by a characteristic-dependent rank test.

V.39.10 From a local gain to a global theorem

Research question V.39.10 (Locate an actual quantitative dependency). Identify a step in a complete progression argument whose bound depends on the local excess $\mathcal{C}_d(F) - \delta^m$ in such a

way that replacing a first-order estimate by the fourth- or sixth-order estimate changes the final quantitative conclusion.

This is the necessary bridge to a stronger Szemerédi-type result. It may turn out that the paper’s dominant losses lie elsewhere, in which case the local sharpening has no effect on the final asymptotic bound. The correct research procedure is to write and audit the full dependency recurrence, not to infer a global improvement from a better constant or exponent in an auxiliary lemma.

[write] **6 October 2026; sources 06–10.** The ten questions are merged into Section V.2 (items III.2, III.16, III.17, III.12, III.10, III.13, III.18, V.2, III.15 and V.1).

V.40 Conclusion

Source 09, delivered Section 15.

Balancing contains more information than a triangle inequality retains. For full cube counts, three exact cancellations move the first error to degree four. The next two homogeneous coefficients admit a complete arithmetic description: parallelograms and stars give ordinary Fourier statistics, while primitive failures of indices two and three give explicit torsion corrections. The multiplicities are simple Stirling-number expressions, and the residual is sixth order with a sharp exponent.

The small-order classification also explains its own limitations. The fifth-order term can be negative; the universal fourth coefficient is attained by a torsion example; and a determinant-five configuration appears one order later. These facts make the refinement more than a formal binomial improvement: they identify exactly which information controls the first deviations from a constant and which arithmetic features a correct generalization must preserve.

The resulting package provides proved local mathematics, exact computational checks, and a concrete route toward formalization. The questions above separate the immediately actionable extensions from the additional work needed for a genuinely stronger global progression bound.

V.41 Proof and artifact dependency map

Source 09, delivered Appendix C.

The logical dependence is as follows. The analytic input is Gowers–Cauchy–Schwarz, the finite Fourier inversion formula, and ordinary finite-dimensional integer linear algebra. Three-vertex primitive uniformity gives the fourth-order bound immediately. The small determinant bounds give the four- and five-vertex classifications; coordinate-block counting gives their multiplicities. Combining these with the subset expansion gives the main exact theorem. Sharpness uses explicit characters and two-frequency functions, with random rounding for indicator transfer. The order-five obstruction has its own determinant and kernel certificate.

The numerical files are independent cross-checks of those mathematical steps. The optional finite table in Appendix III.26 is the one additional exact identity supported by a full explicit enumeration. The source audit and formalization plan do not enter any mathematical proof and do not assert a successful Lean build.

V.42 What the local gains do to a quantitative iteration

Source 10, delivered Section 10.

A better local lemma is useful only after its hypotheses and losses have been propagated through the later argument. We record two exact bookkeeping statements to make that distinction concrete.

Proposition V.42.1 (A reciprocal-density bound on the number of steps). *Suppose $c > 0$, $a \geq 1$, and $0 < \delta_0, \dots, \delta_t \leq 1$ satisfy $\delta_{i+1} \geq \delta_i + c\delta_i^a$ for $0 \leq i < t$. If $a = 1$, then*

$$t \leq \frac{\log(1/\delta_0)}{\log(1+c)}.$$

If $a > 1$, then

$$t \leq \frac{\delta_0^{-(a-1)} - 1}{1 - (1+c)^{-(a-1)}}. \quad (\text{V.42.1})$$

In particular, if $\delta_{i+1} \geq \delta_i + \delta_i^2/320$, then $t \leq 321(\delta_0^{-1} - 1)$.

Proof. For $a = 1$, iteration gives $\delta_t \geq (1+c)^t \delta_0$. For $a > 1$, put $q = a - 1$ and $u = c\delta_i^q \in (0, c]$. Then

$$\delta_i^{-q} - \delta_{i+1}^{-q} \geq c \frac{1 - (1+u)^{-q}}{u}.$$

The function in the fraction is nonincreasing in u , since

$$\frac{1 - (1+u)^{-q}}{u} = q \int_0^1 (1+su)^{-q-1} ds.$$

Thus each decrement is at least $1 - (1+c)^{-q}$. Sum over i and use $\delta_t^{-q} \geq 1$. For $a = 2$, the denominator is $c/(1+c)$; substitution of $c = 1/320$ gives the stated specialization. $\square$

The last estimate sharpens the $640\delta_0^{-1}$ dyadic count displayed in the source's proof of Theorem 2.6, while keeping precisely the same increment hypothesis. It is an elementary refinement of that bookkeeping, not a new form of Roth's theorem. If an argument tries to force a contradiction by performing more steps, it must allow $\lfloor 321(\delta_0^{-1} - 1) \rfloor + 1$ attempted increments, and must verify that the size assumptions remain available at every one.

[write] **6 October 2026; sources 06–10.** The transcription's count $640\delta_0^{-1}$ is in the proof of Theorem 2.6 (exact companion `theorem_2_6_holds`, `Proofs02Roth.lean`). The step counts of source 03's Proposition V.12.1 and source 06's Lemma V.25.1 are compared with this one in the note after source 03's proposition; source 03's (V.12.2) is the next proposition with constant parameters.

Proposition V.42.2 (Exact propagation of length losses). *Suppose $N_{i+1} \geq b_i N_i^{\vartheta_i}$, where $N_i > 0$, $b_i > 0$, and $0 < \vartheta_i \leq 1$. Let $P_t = \prod_{i=0}^{t-1} \vartheta_i$, with $P_0 = 1$. Then*

$$\log N_t \geq P_t \log N_0 + \sum_{j=0}^{t-1} \left(\prod_{i=j+1}^{t-1} \vartheta_i \right) \log b_j.$$

For positive target thresholds $M_0, \dots, M_T$, a sufficient condition for $N_t \geq M_t$ for every $0 \leq t \leq T$ is

$$\log N_0 \geq \max_{0 \leq t \leq T} \frac{1}{P_t} \left[\log M_t - \sum_{j=0}^{t-1} \left(\prod_{i=j+1}^{t-1} \vartheta_i \right) \log b_j \right].$$

For constant $\vartheta \in (0, 1)$, constant $b \in (0, 1]$, and $M \geq 1$, the simpler sufficient condition is

$$\log N_0 \geq \vartheta^{-T} \left(\log M + \frac{\log(1/b)}{1 - \vartheta} \right) - \frac{\log(1/b)}{1 - \vartheta}.$$

Proof. Take logarithms and repeatedly substitute the preceding inequality. This yields the first formula by induction. Solve each of the finitely many resulting lower bounds for $\log N_0$ and take their maximum. For constant parameters the geometric sum evaluates explicitly; its lower bound for $\log N_t$ decreases with t under the displayed condition, so it is enough to check $t = T$. $\square$

For Theorem I.24.2, if the continuous square-root expression is at least two, its floor is at least half that expression. One may therefore take

$$\vartheta_i = \frac{1}{2}, \quad b_i = \frac{1}{2} \sqrt{\frac{\lambda_{\gamma,i}}{2\pi}},$$

provided the appropriate correlation hypothesis is available at stage i . This gives a direct route for comparing choices of the retained correlation fraction γ . It also shows why a claimed improvement in a final tower exponent cannot be inferred from a local constant alone: one must supply the entire sequence of b_i, ϑ_i, M_i and the number of permitted density increments.

V.42.1 The particular limitation in the supplied development

At the pinned snapshot, the edited text places an explicit warning before Theorem 18.1: after replacing an estimate in Corollary 16.11 by the bound supplied by its argument, the displayed assembly no longer establishes the intended quantitative exponent. The source ledger also leaves Corollary 16.11 and Theorems 18.1–18.2 open [48]. These are statements about the repository’s present proof chain. They do not negate the published theorem of Gowers or the subsequent quantitative literature.

The results of this report do not discharge those later dependencies. They can replace specific finite lemmas and reduce some input losses, but a new global bound for $r_k(N)$ would require a separate completed argument through the multilinear extraction, local inverse step, interval transfer, and iteration. The precise propagation formula above identifies what that future argument must certify.

V.43 A map to the Lean development and reproducible checks

Source 10, delivered Section 11.

V.43.1 What has actually been verified

The general results above have mathematical proofs in the article. Independent review checked the noninvertible reindexing, mixed-norm induction, strict tail threshold, exceptional-set denominator, integer partition bounds, diagonal correction in the random selection, and the two representation-counting injections. That review corrected an asymptotic wording issue: when density varies, the length comparison is multiplicative, as stated in Section I.24.

[write] **6 October 2026; sources 06–10.** The independent review mentioned here is not part of the delivery and is not shipped; nothing in this report rests on it (as for source 03’s reviews). The intake reran the four diagnostic scripts on a copy (all exit 0, outputs identical).

The package also contains independent finite diagnostics. Their role is to expose errors in formulas, normalization, and boundary cases; they do not establish a theorem for all finite groups. The executed checks include the following.

Family	Cases	What is checked
Density profiles	1,230	Exact cap and variance inequalities.
Density tails	4,098	The strict-threshold mass estimate.
Exceptional sets	2,532	The retained-cell density guarantee.
Sharp tail examples	45	Exact equality for three-point profiles.
Phase partitions	39,648	Coverage, integer lengths, circle approximation, and 1,408,790 individual cells.
Prime-codomain fibres	23,538	Exhaustive small-field function restrictions, the size guarantee, and the Freiman property.
Weighted representations	41	Exact codegree and representation counts.
Mixed-norm averages	432	Arbitrary complex functions, both AP and general coefficient systems.
Heterogeneous AP counts	210	The weighted telescoping bound.
Torsion sharpness	5	Four characteristic-two dimensions and the exact characteristic-three cube value 49/81.
Cubic norm	125	The determinant identity and unique zero.
Polynomial coloring	390,000	Every ordered nontrivial 4-AP in $\mathbb{F}_5^4$.

The rational profile, fibre, determinant, and explicit coloring checks use exact arithmetic. In the phase script the coverage, circle-step, and half-angle checks are exact rational calculations; complex distances are checked numerically with tolerance 10^{-12} . The arbitrary complex Gowers calculations use NumPy and tolerance 10^{-9} , with random seed 20261006. Their cube powers are computed from the defining sums, independently of the induction used in the proof. The torsion sharp examples are also verified exactly.

The file `verification_results.json` records the actual script outputs and their hashes. The file `source_manifest.json` records the pinned source identities. The README supplies the commands needed to rebuild the PDF and rerun the diagnostics.

V.43.2 Suitable formalization units

The current development defines `UniformOfDegree`, `gowersNorm`, `NatAP`, `ModAP`, `FreimanHom`, and the relevant catalogue statements [48]. Its existing modules suggest the following interfaces. This table is a proposed integration map, not a list of newly compiled declarations.

Mathematical unit	Existing source location	Main obligation
Partition cap and tail	<code>Proofs02</code> <code>DensityIncrement.lean</code>	Finite weighted sums; positivity of $2\delta - S$; equality cases.
Phase partition	<code>NatAP</code> and the Section 2 partition modules	Circle approximation, injective integer progression coordinates, and cardinalities with floors retained.

Mathematical unit	Existing source location	Main obligation
Kernel average identity	Generic finite-group infrastructure, with a cyclic specialization	Exact equal-size fibres of the multiplication endomorphism.
Mixed-norm induction	<code>Proofs03</code> <code>ProgressionAverage.lean</code> and the cube infrastructure	Hölder, derivative moments, and the precise invertible-neighbor hypothesis.
Weighted telescoping	<code>Proofs03</code> <code>ProgressionCount.lean</code>	Vanishing of the first two terms and a separate count of repeated progressions.
Weighted energy extraction	<code>Proofs07BalogSzemerédi.lean</code>	Finite random sampling or equivalent counting, diagonal correction, and injectivity of the tuple encodings.
Obstruction fibre and torsion budget	<code>Proofs07</code> <code>FreimanRestriction.lean</code>	The graph projection injection; exact bad-step fibres; translation averaging.
Polynomial obstruction	A new finite-field polynomial module	Coefficient homogeneity, a Vandermonde determinant, Chevalley–Warning, and the norm construction.

For the mixed-norm theorem, it may be easier first to formalize its equivalent version after raising to the power 2^{k-1} :

$$|\Lambda_{\mathbf{a}}(f_1, \dots, f_k)|^{2^{k-1}} \leq K \|f_k\|_{U^{k-1}}^{2^{k-1}} \left(\prod_{i=1}^{k-2} (\mathbb{E}|f_i|^{2^i})^{2^{k-1-i}} \right) \cdot (\mathbb{E}|f_{k-1}|^{2^{k-2}})^2.$$

This formulation avoids fractional powers in the inductive algebra. Taking nonnegative roots can be a final wrapper.

The probabilistic language in Theorem II.20.1 also has an entirely finite interpretation. For a fixed sampled $v \in A$, sample independent integers $u_a \in \{1, \dots, m\}$ and retain a when $u_a \leq r_A(v - a)$. This exactly realizes the required rational membership probabilities. The proof can therefore be phrased using finite sums over $A \times \{1, \dots, m\}^A$, without introducing a general probability-space interface.

The most important source-to-theorem checks remain mathematical: the normalization must match, positivity must justify each reindexing bound, the progression must have the claimed distinctness, and the new theorem’s assumptions must actually be available at its call sites. A successful axiom audit alone would not verify those comparisons.

V.44 Further research questions and proposed next steps

Source 10, delivered Section 12.

The questions below separate proposed refinements of the present theorems from established questions in the wider literature. A question introduced here is not asserted to be previously unknown or to have been recognized as open by the community.

V.44.1 Finite problems closest to the proved results

1. **Optimal torsion constants beyond the sharp families.** For a fixed finite abelian G , coefficient pattern, and target having at least one invertible coefficient difference, determine the least constant C in $|\Lambda_{\mathbf{a}}(f)| \leq C \|f_j\|_{U^{k-1}}$ for 1-bounded functions. Theorem III.28.2 gives an explicit upper bound, and Proposition III.28.4 settles particular families. The next small case is the exact constant for four-term patterns on elementary abelian 3-groups. Small finite optimizations can suggest extremizers, but an analytic extremal proof is the target.
2. **Use restricted direction averages before losing them.** The fibre identity retains the more informative quantity $\mathbb{E}_h \mathbf{1}_{cG}(h)V(h)$. Determine when the images cG arising at successive Cauchy–Schwarz steps have enough compatibility to improve the product of kernel cardinalities. An exact comparison for nested images would already strengthen the generic bound.
3. **Optimize the geometric part of the phase tradeoff.** The density cap is sharp, but the passage from one approximate period to blocks of lengths L through $2L - 1$ may leave room in the constants. Seek the best universal minimum block length for a prescribed phase disc radius, with all integer rounding retained. A useful first goal is a matching obstruction or a better construction at fixed η/δ as $\delta \rightarrow 0$.
4. **Simultaneous correlation data.** If several approximately constant characters or polynomial phases are available on the same partition, their correlations can force a larger conditional second moment than one correlation alone. Determine the best Gram-matrix bound for that moment and combine it with (I.23.3). The geometry must supply the common partition; assuming it for free would conceal the main cost.
5. **A sharp vertical-fibre restriction theorem.** How close is $|B|/[k(h_k - 1) + 1]$ to the best bound determined only by the defect fibre? Can the factor k be reduced under additional additive structure of H_k ? The exact case $H_k = \{0\}$ is settled, but the extremal nonzero-fibre case remains a natural proposed problem.
6. **Exploit overlaps in the bad-step union.** Theorem II.21.3 uses a union bound over $E_j = \{d \neq 0 : jd \in F\}$. Exact intersections between these sets can be substantial. A formula adapted to a cyclic prime-power codomain, using valuations and the distribution of defects in jV , is a concrete route to a better general-codomain retention bound.
7. **Optimize the graph-energy input and its propagation.** The explicit proof here retains a proportion $\gamma/4$ before Freiman extraction; the published Reiher–Schoen theorem retains almost $\sqrt{\gamma}$ [12]. Propagate the latter, with the exact vertical fibre when available, through the later Bohr-neighborhood and local-linearity stages of the supplied proof. The desired output is a complete new chain of numerical inequalities, not an isolated replacement of the exponent 1164.

V.44.2 Polynomial colorings after the exact low-degree answer

8. **Higher degrees.** Theorem III.30.3 completely answers the specified coordinate problem for $D_\ell \leq k$. Above that range, the higher-degree obstruction proved in Section III.30 gives necessary dimension bounds without matching constructions. For $k = 4$, a scalar polynomial of degree five or six has the resulting upper dimension bound eight. Determine the largest

attainable dimension, starting with dimension five. This is a proposed finite extremal problem, not a claimed established open conjecture.

9. **Stability at the dimension threshold.** At $d = \sum_{\ell} \lceil D_{\ell}/2 \rceil^2$, must every avoiding polynomial have some structural feature resembling an anisotropic odd-degree decomposition? Full equivalence to the norm construction may be too strong, but a classification in the first cubic four-dimensional case would test the idea. Coordinate changes and linear changes of the color coordinates should be factored out before comparing examples.
10. **Actual colors versus declared coordinates.** The exact theorem controls the number of polynomial coordinates. Can a map with more coordinates but an unusually small image outperform the full-palette exponent? A useful result would relate avoidance, fiber geometry, and actual image cardinality. The theorem proved here does not supply such a relation.
11. **Unrestricted symmetric-progression colorings.** Deng, Tidor, and Zhao ask whether $N^{o(1)}$ colors suffice to avoid nontrivial symmetrically colored four-term progressions [37]. Shi and Dong’s $O(N^{1/4})$ construction improves the known polynomial exponent [33]. The dimension obstruction here explains a limit of one bounded-degree polynomial-coordinate model. It leaves unrestricted colorings, higher-degree constructions, and more economical treatment of digit information as meaningful directions for further work.

V.44.3 A practical order of work for this repository

The first implementation target should be the optimal partition cap, tail bound, and kernel-fibre averaging identity: they are small, independent facts with clear finite-sum statements. Next come the mixed-norm progression theorem and weighted counting wrapper, where normalization and torsion hypotheses are the main integration risks. The obstruction-fibre injection is another compact structural target.

The weighted energy proof is a larger, but still self-contained, formalization project. Its exact finite sampling and representation injections make its dependencies transparent. The polynomial dimension theorem then forms a separate finite-field development with a complete mathematical endpoint and a useful explicit example.

Finally, any effort to improve the headline Szemerédi estimate should return to the unresolved quantitative dependencies in the source ledger, especially the transition through Section 16. Once those hypotheses and bounds are discharged, the local improvements in this report can be inserted into an audited parameter calculation using Propositions V.42.1 and V.42.2.

[write] **6 October 2026; sources 06–10.** The eleven numbered questions, and the two question paragraphs at the ends of Sections III.29 and II.21, are merged into Section V.2 with credits. The intake dossier counted ten numbered questions for this source; there are eleven (seven in its first subsection, four in its second).

V.45 External inputs, provenance, and reproducibility

Source 10, delivered Appendix A.

The new quantitative deductions rely on elementary finite algebra and analysis except where the following established inputs are explicitly used: Fourier orthogonality and Parseval, finite-field extension norms, and the Plünnecke–Ruzsa sum-difference inequality. The needed Chevalley–Warning statement is proved in the article. The Reiher–Schoen result is used only for the separately

labeled modern-input alternative; the explicit $2^{-209}\gamma^{103}$ corollary uses the energy theorem proved here.

The source comparison was made against the pinned repository files, including the paper transcription, `Definitions.lean`, the numbered Sections 1–3 and 6–7, the relevant density/counting modules, the Section 18 quantitative wrapper, the status ledger, and the axiom audit. The manifest gives exact paths and blob identifiers for these materials. The bibliography points to the original publication separately from the edited repository version.

The main $\text{\LaTeX}$ file contains the complete article and bibliography. Its only external illustration is `figures/phase_frontier.pdf`; the script `verification/make_figure.py` regenerates it and the numerical examples. The plotted points evaluate theorem guarantees and do not represent observed extremal sets. The four mathematical diagnostic scripts can be run separately or through `verification/run_all.py`. None requires network access.

V.46 Finite verification and formalization interfaces

Source 11, delivered Section 12.

V.46.1 What was computed

The accompanying standard-library Python program uses integer arithmetic for every energy and inequality check. The main exhaustive computation subtracts the unique affine map interpolating $0, e_1, \dots, e_n$. Every map then has a unique representative with zero values at those points. Since affine subtraction preserves energy and distance to the affine class, enumerating these normal forms is an exhaustive computation modulo a proved symmetry, not a sample.

Domain and target	Normal forms	Maps represented	Maximum	Extremizers
$\mathbb{F}_2^2 \rightarrow \mathbb{F}_2$	2	16	40	8
$\mathbb{F}_2^3 \rightarrow \mathbb{F}_2$	16	256	344	128
$\mathbb{F}_3 \rightarrow \mathbb{F}_3$	3	27	15	18
$\mathbb{F}_5 \rightarrow \mathbb{F}_5$	125	3125	69	500
$\mathbb{F}_7 \rightarrow \mathbb{F}_7$	16807	823543	211	2058
$\mathbb{F}_3^2 \rightarrow \mathbb{F}_3$	729	19683	489	486

Here “maximum” and “extremizers” refer to non-affine maps. Every computed extremizer was checked to have distance exactly one from the affine class. The six runs represent 846650 maps through 17682 normal forms. Multiple representations of a one-point perturbation by different affine maps in the very smallest cases do not affect the table: actual functions, not descriptions, are counted.

Additional checks comprise all 2176 nonzero labelings of two-point supports on $\mathbb{Z}/17\mathbb{Z}$ into $\mathbb{Z}/5\mathbb{Z}$, 1800 seeded odd-target sparse examples, and 1000 seeded examples with targets of even order, including mixed targets such as $\mathbb{Z}/6\mathbb{Z}$. These verify the stronger comparison and deficit formulas, not just the leading cubics. Three subgroup equality examples and an explicit two-torsion counterexample to the odd formula are also included. The first eight inverse coefficients are checked by exact substitution.

The files `checks/verification_results.json` and `checks/verification_log.txt` record the output. Finite checks do not prove universal statements, do not establish literature priority, and are not Lean kernel checks.

V.46.2 The inspected repository snapshot

The relevant source was inspected at ProveIt’s commit

40f885098a74607b423909dbcb8f070993b7c2a8.

The inspected files have the following blob identifiers [60, 61, 62]:

Paper transcription	41085a0efde8c1932e86e80c791984f52841c22e
Sections06_07.lean	57b79242466ca9ddf5c36c76b828d7c5557f58ca
Proofs05_10.lean	6a49bbe6a307e6a4a20ca5d014e173fcd20fc22d

The statement file explicitly distinguishes proposition-valued definitions from proofs. It defines `functionGraph`, `proposition_6_1`, and `lemma_6_2`; the latter uses `Finset.addEnergy`. The existing convention represents a partial map by a total map together with a finite domain. That convention is compatible with Theorem II.32.1.

V.46.3 A practical formalization decomposition

The proposed extension should preserve this interface rather than reinterpreting the normalizations. A useful dependency order is:

1. Prove the derivative-collision identity for a function graph over finite additive commutative groups, and affine-subtraction invariance. Keep cardinalities in $\mathbb{N}$ before introducing normalized real quantities.
2. Formalize the six support-pattern classes and their exact counts. The only changes of variables needed are bijections between G^3 and additive quadruples, so no primality hypothesis belongs here.
3. Introduce nonzero error fibers, M, A, C, J , and prove the two color inequalities. Make the no-two-torsion assumption explicit in $C \leq A$; it is not needed for $C \leq M$.
4. Formalize the two comparison inequalities and their integer-valued master bounds. Derive the real cubic inequalities only afterwards. The coset equality case reduces to the elementary equality case of $E(S) \leq |S|^3$.
5. Formalize the anchor average and majority correction using finite probability or normalized finite sums. Then prove the unique-root comparison without importing an analytic power-series theory.
6. Add the partial-domain corollary and instantiate the result after `proposition_6_1`. The vector-space quotient proof can be a separate file, avoiding unnecessary linear algebra in the general-group stability theorem.

No Lean executable was available in the working environment, and no Lean build was run for these new results. The package therefore includes a formalization plan, not a purported verified patch. No repository files were modified. Statements read in the repository and computations run locally have been kept separate from claims of kernel verification.

V.47 Further research questions

Source 11, delivered Section 13.

The following are proposals arising from the proved results. They are not claims that each formulation has previously appeared as an open problem.

Research question V.47.1 (The full small-distance energy envelope). For fixed G, W and an attainable affine distance t , determine

$$\max\{\mathcal{E}(f)/|G|^3 : d(f) = t\}.$$

How far can the odd-target radius $1/7$ and the involution-target radius $1/3$ be enlarged? The current comparison proof loses information when three-error gains are bounded independently of two-error losses. A sharp joint optimization of these patterns may improve the radii without changing the cubic near zero.

Research question V.47.2 (The next energy levels). After excluding affine maps and one-point perturbations, determine the next largest graph energies and all extremizers, first for $\mathbb{F}_p \rightarrow \mathbb{F}_p$. Two-point supports allow constant, opposite, and unrelated error values. The exact support-pattern formulas suggest a manageable finite classification, but a proof must also exclude functions farther from every affine map.

Research question V.47.3 (A complete near-extremizer theorem at the right scale). Combine (II.27.1)–(II.27.3) with a quantitative coset-stability theorem for $E(S) \geq (1 - \theta)|S|^3$. Seek bounds in relative symmetric difference that remain meaningful when $t \rightarrow 0$. The natural parameter supplied by the current argument is η/t^3 , not just η .

[write] **6 October 2026; sources 11–14.** *Answered in substance* by composing Corollary II.27.1 with source 14’s Theorem II.43.2: Remark W3 gives, with $\theta = \eta/(6t^3) \leq 1/100$, a unique subgroup coset C with $|S \triangle C| \leq \tau(\theta)|S|$ and a single error value off a further $\eta/(2t^2(1 - 7t))$ fraction of S , in exactly the parameter η/t^3 asked for; it is written in neither source. Item II.20.

Research question V.47.4 (Small arbitrary groups and mixed torsion). Remove the sufficient condition $|G| \geq 34$ in the unrestricted-group theorem, or classify its exceptions. The vector-space proof uses the ability to extend a homomorphism from the constant-derivative subgroup; that extension can fail for general finite abelian groups. Understanding this obstruction may be as important as refining numerical constants.

[write] **6 October 2026; sources 11–14.** Remark W2 records that the cutoff cannot be removed for general groups $(\mathbb{Z}/12\mathbb{Z} \rightarrow \mathbb{Z}/2\mathbb{Z})$: 1296 against 1288) and the small exceptions found by exhaustive search. A possible lowering of the cutoff to 30 through source 12’s majority lemma is a marked question in item II.21.

Research question V.47.5 (Separate budgets for erasures and corrupted values). The partial-domain theorem treats holes through completion and a coarse cardinality inequality. Determine a sharp two-parameter envelope when a prescribed proportion of inputs is absent and a prescribed proportion of the remaining values is corrupted. Deleting a point and changing its value have different second-order energy costs, so one parameter cannot capture all the available information.

Research question V.47.6 (Efficient certified recovery). Can the canonical affine map be recovered using nearly linear time in the input-table length, with an a posteriori certificate matching the cubic bound? Randomized derivative sampling may reduce the anchor cost, but uniform recovery of the modal homomorphism must be justified, and the rare-error regime requires explicit sample-complexity accounting.

Research question V.47.7 (Higher-order respected arrangements). Develop analogues for the multilinear arrangement counts used later in Gowers’s paper. The key issue is not merely whether a higher-order defect is small: the pattern of repeated coordinates and the torsion-dependent cancellations must be tracked well enough to identify the true extremal error geometry. A naive replacement of 4 by the number of cube vertices would not settle the equality problem.

Research question V.47.8 (Can a high-energy endpoint be reached inside the general proof?). The present refinement is strongest when the respected-quadruple density is near one, whereas the initial inverse step generally gives much less. Is there a localization or amplification procedure that reaches this endpoint without losing more quantitative information than it gains? A useful answer must provide the size of the retained domain and the energy normalization together; replacing N^3 by a small-set cube without an extension argument would not suffice.

Research question V.47.9 (A machine-checked extremal layer). Implement the support-pattern counting and color-comparison layer as a reusable Mathlib development, then connect it to the existing ProveIt Section 6 statements. The finite certificates included here are useful regression tests, but a kernel-checked proof should make the two-torsion split and every normalization explicit.

V.48 Conclusion

Source 11, delivered Section 14.

The exact-additivity endpoint in Gowers’s Section 6 admits a precise quantitative theory. The principal mechanism is an error-value comparison: for sparse enough errors, replacing the nonzero values by one suitable value can only increase graph energy. Exact support-pattern counts then produce a cubic, and the equality case of set energy identifies subgroup cosets. Majority correction turns this local statement into a high-energy inverse theorem. The integer spacing of Hamming distance yields an exact non-affine energy gap; a quotient-by-affine-directions argument handles all prime-field vector-space sizes.

The results also delimit what this approach has not accomplished. They do not improve the small-correlation exponent in Proposition 6.1, do not supply a new global Szemerédi bound, and are not yet Lean-verified. What is supplied is a self-contained collection of sharp finite statements, explicit torsion corrections, equality classifications, reproducible checks, and a concrete route for extending the formalized Section 6 interface.

V.49 Reproduction and provenance

Source 11, delivered Appendix A.

The package contains the article source and compiled PDF, the exact verification program and its recorded output, a build script, and a short formalization interface note. A local reproduction requires a standard LaTeX installation with the packages used in the preamble and Python 3. From the package directory run

```
python3 checks/verify.py --output checks/verification_results.json
pdflatex -interaction=nonstopmode -halt-on-error article.tex
pdflatex -interaction=nonstopmode -halt-on-error article.tex
```

A third LaTeX pass can be used if cross-reference warnings remain after a source edit. The verification program uses only the Python standard library and deliberately avoids floating-point inequality tests.

The repository reads were limited to the cited paper interface and relevant Lean statement/proof excerpts. They are not a repo-wide audit. The bibliography records the source locations used for context and comparison. Original source PDFs and third-party font files are not redistributed in this package.

V.50 Formalization interfaces and verification

Source 12, delivered Section 10.

V.50.1 Where the new statements fit

The relevant definitions in the repository are `bohr`, `FreimanHom`, `Section10Setup`, and `HasBohrDifferenceModel`. The first uses centred modular absolute values and closed inequalities; the second is based on Mathlib’s additive Freiman-homomorphism predicate. The last two package the quantitative input and the local output of Section 10 [64].

New interface	Existing source and implementation role
Finite translation-loss and tree-intersection lemmas	Generalize the counting helpers in <code>Proofs10Bohr.lean</code> ; need only finite-set cardinalities and translation bijections.
Canonical difference and sharp extension	Replace the four-translate witness used for <code>corollary_10_11_holds</code> by a three-vertex argument and an induced value on $C + C$.
Finite radius selection	Add a covering lemma and a product-of-ratios argument; no differentiation, integration, or full regularity theorem is needed.
Fourth-moment smoothing	Reuse Parseval and the fibre cap in <code>Proofs10Main.lean</code> ; avoid the separate cardinality bound on the large-spectrum fourth-power sum.
General synchronization	Parameterize the overlap constant in <code>Proofs10Induced.lean</code> ; retain the nonempty-fibre hypothesis.
Parameterized local-model theorem	Separate <code>theorem_10_13_holds</code> from its fixed density schedule, choose B before the component, and retain the exact $\varrho^2\alpha$ slice fraction.

Table V.50.1: Proposed interfaces. These are a formalization plan, not declarations already checked by Lean. All named files are in the repository’s `GowersSzemerédi` directory.

An economical order is to formalize lemmas II.34.3 and II.39.1 and theorems II.35.2, II.37.2 and II.38.1 independently, then assemble theorem II.40.2. In particular, the extension lemma should not be tied to Bohr sets. This permits it to serve approximate groups, finite intervals, and other approximately translation-invariant domains.

The numerical schedule should be a separate theorem whose conclusion is the conjunction of (II.40.3)–(II.40.4), the rank bound, and the final radius comparison. This makes a change in

constants locally reviewable and prevents a lost ceiling or missing density hypothesis from being concealed in a large proof.

V.50.2 Independent finite checks

The accompanying scripts provide reproducible consistency checks. They supplement the proofs and do not certify arbitrary groups or constitute Lean verification.

Check	Cases	Arithmetic and scope
Three-translate lower bound	560,760	Exact integers; all nested $S \subseteq B$ and ordered directions in cyclic groups of orders at most 8.
Sharp-boundary construction	90	Exact fractions; $1 \leq m \leq 12$ and every $0 \leq t \leq m$.
Selected Bohr radius	7,900	Exact rational radii in 12 cyclic and product groups.
Overlap of the selected set	16,932	Exact set counts, including nonzero shifts which cross a nonempty boundary.
Fourier smoothing estimates	17,970	Direct finite convolutions and numerical Fourier transforms; maximum fourth-moment identity discrepancy $3.33 \cdot 10^{-16}$.
Directional repair examples	4	Exact pass/fail inequalities, including missing points and corrupted values in cyclic groups of orders 31, 47, and 101.

Table V.50.2: Executed checks included in the archive. The numerical Fourier tolerance is not used in any mathematical proof.

For example, the full-group repair test on $\mathbb{Z}/31\mathbb{Z}$ changes one value of an affine map. Its largest directional collision defect is 118/961. The majority map recovers the original slope, and the repaired affine map agrees at 30/31 of the points. A second test deletes three points and changes two values; the exact collision criterion still guarantees the recovered model, which agrees on 13/14 of the remaining set. The local interval test deliberately makes no global-affine claim.

V.50.3 Status of the deliverable

The archive contains the complete $\text{\LaTeX}$ source, compiled PDF, the figure and its generating source, both verification scripts, their executed reports, and a short build/readme file. The article's propositions and proofs are mathematical derivations; the supplied code checks finite instances. The inspected repository is identified by an immutable commit, and no change to that repository is part of this deliverable.

V.51 Further research questions and topics

Source 12, delivered Section 11.

The questions below are proposed research directions arising from the proved results. Except where an external theorem is cited, they are formulations for further investigation, not assertions that the exact wording is a previously published open problem. They are ordered roughly from local structural questions to the full quantitative program.

Question V.51.1 (Classify extremizers for the sharp overlap tradeoff). *For which B, S, C, ϕ does extension fail with $3\varepsilon + 2\kappa = 1$? Are irreducible equality examples governed by a three-part quotient, as in theorem II.35.3, or are there fundamentally different configurations?*

The proof indicates a concrete approach: equality requires the missing translates and the two boundary losses in a failed triangle to exhaust the ambient set without waste. Classifying these equality cases would be more informative than adjusting a numerical constant.

Question V.51.2 (Stability near the failure boundary). *If extension fails and $3\varepsilon + 2\kappa \leq 1 + u$ with small $u > 0$, what structural approximation to an extremal example follows under natural nondegeneracy assumptions?*

The assumptions matter: one can add unused components or directions to destroy a naive classification. A useful statement should specify how much of B and of the direction set participates in failed relations, then control the symmetric-difference distance to a model configuration.

Question V.51.3 (Prime cyclic groups and the same target). *Determine the optimal deterministic density-overlap region when $G = H = \mathbb{Z}/p\mathbb{Z}$, especially when B and C are Bohr sets with the same spectrum. How much does this narrower class improve on the universal $3\varepsilon + 2\kappa < 1$ guarantee?*

The construction proving universal sharpness does not settle this case. A study should first compare existing dense-set rectification and critical-pair results, then use finite linear algebra to identify the smallest prime-field obstructions. The distinction between a theorem for all abelian targets and one for a fixed target is essential.

Question V.51.4 (Optimal constants for finite radius selection). *For $2/3 < a \leq 1$, what is the best universal $c(a)$ such that some $r \in [\rho/2, \rho]$ supports extension from every a -dense subset of $\mathcal{B}(\Gamma, r)$ to $\mathcal{B}(\Gamma, c(a)\rho/d)$?*

The present argument gives $c(7/8) = 1/8$. It passes through the covering factor 4^d and through an inner-set lower bound for the overlap. Direct overlap estimates can be stronger than the inner-set count. Product examples show the $1/d$ scale is natural, but leave the constant unresolved.

Question V.51.5 (Measured growth in place of nominal rank). *Can a later stage of Gowers's argument preserve and use the actual growth ratio $|\mathcal{B}(\Gamma, \rho)|/|\mathcal{B}(\Gamma, \rho/2)|$, instead of replacing it by 4^d ?*

Theorem II.37.2 already gives the corresponding local improvement. The next issue is whether the rank or growth data can be carried through simultaneous families of spectra without losing the relations that motivated Section 10.

Question V.51.6 (Anisotropic spectral neighbourhoods). *Can separate radii for different characters exploit the weights $|\widehat{g}(\gamma)|^4$ while retaining the structural compatibility required later in the proof?*

The large-spectrum error before taking a common radius is bounded by

$$\frac{2\pi}{N} \sum_{\gamma \in \Gamma} \rho_{\gamma} |\widehat{g}(\gamma)|^4.$$

This suggests a constrained optimization problem for the radii. Any gain must pay for the resulting volume and for the cost of common refinements; maximizing the volume of one isolated box is not sufficient.

Question V.51.7 (From average error to uniform directional repair). *Under what additional overlap or representation hypotheses does a small global proportion of bad additive quadruples force the collision hypothesis of theorem II.36.2 on a useful Bohr neighbourhood?*

An answer could shorten the component-extraction part of Section 10. The challenge is to control directions with few representations. The weighted budget after theorem II.36.2 suggests retaining actual edge counts rather than discarding all information through a uniform worst-case bound.

Question V.51.8 (Sharp robust correction bounds). *Determine the best repair region in $(\varepsilon, \kappa, \eta)$, and the optimal affine-agreement conclusion in the global case. Is the nested bound $1 - \chi(\chi(\eta))$ attained subject to compatibility between all directional distributions?*

Each separate collision-to-majority inequality is sharp, but their simultaneous attainability in one group-valued function is an additional constraint. The factor 3 in the error budget also comes from a union bound which may not be extremal for compatible increments.

Question V.51.9 (Higher-order local relations). *Is $3\varepsilon + r\kappa < 1$ the correct universal scale for forcing an order- r Freiman map on C from an order-two map on S ? Which spanning-tree or relation-complex parameters replace r for a prescribed finite family of additive identities?*

The canonical difference map allows long relations to be checked through short triangles, so the density loss need not grow with the number of summands. More efficient decompositions of the prescribed relations may improve the boundary term.

Question V.51.10 (Reassemble Sections 13 and 16 with explicit new budgets). *Reprove the later local-model inputs using theorem II.40.2, retaining both the polynomial radius and the exact slice fraction $(1 - q)\varrho^2\alpha/16$.*

This is the most direct continuation of the present article. It requires a dependency-by-dependency calculation, including spectra indexed by other variables, common progressions, multilinear compatibility, and each minimum-width condition. The current repository's unresolved Corollary 16.11 estimate should be resolved on its own terms before a global numerical conclusion is drawn [66].

Question V.51.11 (A verified parameter calculus for the whole proof). *Can the local theorems expose monotone, composable bounds on rank, radius, retained density, and minimum ambient size so that the final iteration is a separately verified calculation?*

The parameterized theorem here is a useful first interface. A complete implementation should distinguish a mathematical hypothesis from a bound introduced merely for convenience, and should retain exact cardinality factors until the final simplification. This would make future improvements easier to assess and prevent local gains from being lost invisibly.

Question V.51.12 (Measure and infinite-group analogues). *Extend the translation-loss and synchronization arguments to finite-measure subsets of locally compact abelian groups, with Haar measure replacing cardinality and appropriate measurability assumptions on the maps.*

The finite intersection inequalities are measure inequalities, so much of the algebra survives. The points requiring care are representative selection, almost-everywhere versus everywhere identities, continuity of an induced homomorphism, and the meaning of a positive-measure family of witnesses. These should be stated explicitly rather than inferred from the finite case.

V.52 Parameter ledger and reproducibility

Source 12, delivered Appendix B.

The following ledger lists the fixed-spectrum specialization without dropping a density factor.

Symbol	Value or invariant
η	2^{-43} , error in the order-eight hypothesis.
θ, q	$10\eta^{1/5}$ and $\sqrt{\theta}$, with $\theta < 1/36$, $q < 1/6$.
α	$ X /(MN)$, with $0 < \alpha \leq 1/6$ in the fixed-spectrum specialization.
ϱ	$\alpha^2/32$, used for component selection, not a Bohr radius.
σ	$2^{-72}\alpha^9$, relative invariance error for fourfold fibres.
λ	$2^{-37}\alpha^{11/2}$, unchanged spectral threshold.
d	Actual cardinality of Γ ; $1 \leq d \leq 2^{74}\alpha^{-10}$.
ρ	$2^{-74}\alpha^9/\pi$, outer Bohr radius furnished by the fourth moment.
r_0	One of $3d$ selected radii in $(\rho/2, \rho]$, chosen before S and W .
κ	$1 - 4^{-1/3} < 3/8$, guaranteed translation loss on the selected B .
L	$M X ^3$, the fourfold-domain fibre cap.
$ \Omega $	$ X ^4 = \alpha LN$; the effective density is still α .
$ W $	At least $\varrho^2\alpha^2LN/16$.
$ W_1 / \Omega $	At least $(1 - q)\varrho^2\alpha/16$.
$ Y / X $	The same fraction after fixing three coordinates; at least $\alpha^5/20000$.
ζ	$\rho/(6d)$, at least $\alpha^{19}/(3\pi 2^{149})$.

The source snapshot was inspected through the GitHub connection. The source texts are cited by immutable commit in the bibliography. The final verification reports in the archive are produced by:

```
python3 verify_extension.py
python3 verify_bohr_fourier.py
python3 make_figure.py
pdflatex -interaction=nonstopmode -halt-on-error gowers_bohr_extension.tex
pdflatex -interaction=nonstopmode -halt-on-error gowers_bohr_extension.tex
```

The first verification uses the Python standard library. The second and the figure generator use NumPy and Matplotlib as indicated in the readme. The article does not require external bibliography software or network access to compile.

V.53 Finite examples, verification, and formalization

Source 13, delivered Section 10.

V.53.1 A complete partition beside a sharp obstruction

Take $N = 1009$, $r = 1000$, and the two modular phases

$$\phi_1(n) = 92n, \quad \phi_2(n) = 868n \pmod{1009}, \quad \varepsilon_1 = \varepsilon_2 = \frac{1}{4}.$$

The longest simultaneous good progression has exactly seven elements. For length eight, a candidate common difference is at most $\lfloor 999/7 \rfloor = 142$. The no-wrap criterion would require both centered steps to be at most $1009/28$. Checking all 142 possible differences shows that none qualifies. Length seven is attained with difference 143: its centered phase steps are 39 and 17, and the two spans are 234 and 102, both at most $1009/4$.

The same interval nevertheless has a complete common-difference partition into 330 cells of lengths three or four, with difference 22, produced by theorem I.26.8. This example illustrates why complete coverage and one long progression have related but distinct finite optimization problems. Its data, including the complete partition construction and all difference checks, are reproduced by `verification/verify_partitions.py`. The finite union-bound certificate also holds: its left-side count in (I.26.3) is 756718, below $1009^2 = 1018081$.




Figure V.53.1: The exponent comparison for two phases. The flexible boundary is optimal in the growing-cell sense proved in corollary I.26.6. The prescribed-length line is the sufficient boundary from the conductor argument; its optimality is not asserted. The line $b = 2a$ is the simultaneous length/diameter comparison used for Lemma 5.9.

V.53.2 What the executable checks establish

The accompanying scripts are deterministic and use no network. They address fragile parts of the derivations—normalization, integer rounding, orbit counts, Fourier signs, and equality cases. They support reproducibility; the mathematical proofs above establish the general theorems.

Script	Checked instances	Main checks
<code>verify_partitions.py</code>	3,437 partitions; 579,417 points; 410 composite-modulus cases	Exact coverage, disjointness, declared lengths, modular diameters, conductor thresholds; an explicit seven-point obstruction example
<code>verify_density.py</code>	50,000 seeded mixtures; five equality constructions; 440 parameter cases	Cell envelope, aggregate envelope, joint frontier, exceptional mass, Fourier-increment endpoints and rounding
<code>verify_cubes.py</code>	All 256 subsets over $\mathbb{Q}$ and $\mathbb{F}_p$, $p = 3, 5, 7, 11, 13$; 15 function cases	Rational cancellation and orbit identities; complete centered expansion; face-square identities; Fourier formulas and sign examples
<code>verify_fourier_counting.py</code>	8,146 intervals; 132 AP families; 186 spectral cases	Finite interval constant, exact dual equality, telescoping, endpoint uniformity bounds, a negative control, spectral-deficit equality

The partition tests use exact integers and rational arithmetic. Cube combinatorial and physical-space identities use exact rational arithmetic; Fourier evaluations use floating point, with maximum observed discrepancy about $3.5 \cdot 10^{-18}$ in that suite. The interval dual equality error was below $7 \cdot 10^{-16}$, and the telescoping identity error below $9 \cdot 10^{-17}$. The random density tests allow 10^{-12} tolerance, explicitly recorded in their output. Counts, seeds, versions, examples, and errors are stored in the accompanying JSON files. These tolerance checks are not described as exact arithmetic proofs.

V.53.3 How to fit the results into the Lean development

The existing source distinguishes proposition-valued numbered statements from companion theorems proving them. A robust extension should preserve those named interfaces and add stronger reusable lemmas underneath. The following table gives concrete attachment points in the inspected snapshot [68, 69].

Inspected interface	Suggested stronger helper	Main formal work
<code>Section05.lean, lemma_5_9</code>	Simultaneous affine partition with explicit target lengths	Rectangular finite pigeonhole; residue chains; consecutive-length decomposition; diameter transport
<code>Section05.lean, lemma_5_15</code>	Weighted high-average mass and discarded-mass bounds	Positive/negative parts; finite weighted sums; denominator positivity; extremal construction
<code>Proofs05 Downstream.lean, local phase refinements</code>	Squared real-phase envelope and joint secant bound	Complex real-part inequality; endpoint envelope; convex norm of an affine map
<code>Sections01_03.lean, corollary_2_5</code>	Density-sensitive Fourier increment	Compose the new one-phase partition, phase removal, and weighted maximum-average lemma

Inspected interface	Suggested stronger helper	Main formal work
<code>Proofs03CubeUpper.lean</code> , <code>lemma_3_10_holds</code>	Coloop cancellation and exact centered expansion	Finite field linear algebra; measure-preserving translations; subset enumeration; normalized cube forms
<code>Proofs03Interval.lean</code> , <code>lemma_3_5_holds</code>	Length-sensitive interval spectral bound	Retain the minimum of length and decay; finite/integral comparison; normalize Fourier coefficients
<code>Proofs03</code> <code>ProgressionCount.lean</code> , <code>corollary_3_3_holds</code>	Telescoping count with density weights	Algebraic telescope; exact two-form factorization; reuse endpoint generalized von Neumann

The recommended initial formalization target is theorem I.27.1: it uses only finite weighted real sums and exposes a useful generic API. Next come the two finite integer partition decompositions and the one-phase Fourier corollary. The full three-cube identity is longer but has a finite, auditable structure: certify the coefficient-matrix kernels, prove the orbit identifications, and rewrite the resulting sums.

Three distinctions should remain explicit in any implementation. First, the article’s norms are normalized while the source `gowersNorm` is not. Second, cardinality-equals-length is not the same as an integer nonwrapping lift. Third, an axiom audit validates the dependency boundary of compiled theorems, but does not discharge hypotheses of conditional theorems or verify the correspondence between a proposition and its intended paper statement. The repository’s audit comments make this distinction themselves. No new Lean source with an unverified compilation claim is included in this package.

V.54 Further research questions and concrete next steps

Source 13, delivered Section 11.

The following questions are suggested by the proved results. Their status here is deliberately local: the report does not claim that each formulation has never appeared or been resolved elsewhere. Before treating one as a historical open problem, its exact formulation should receive a dedicated literature check.

V.54.1 Complete partitions with prescribed lengths

Research question V.54.1. Can a complete partition into lengths $v - 1$ or v achieve the flexible exponent frontier $(q + 1)a + qb = 1$, or is there a genuine arithmetic obstruction between that frontier and the proved sufficient region $(q + 2)a + qb \leq 1$?

[write] **6 October 2026; sources 21–27.** *Answered by source 23:* there is a genuine arithmetic obstruction. The prescribed-length frontier is $(q + 2)a + qb = 1$, even with cell-dependent differences (Proposition I.38.1), and the threshold is $L^{q+2} / \prod_i \varepsilon_i (1 + o(1))$ for lengths $L, L + 1$ (Corollary I.37.1); the lower example is realizable on a prime cyclic group. With an exceptional set the cost is quantified by its Theorem I.39.1 and Corollary I.39.4.

The present proof uses one common difference and requires every residue chain to pass the conductor $(v-1)(v-2)$. A useful first step is to formulate three separate minimax problems: one common difference, cell-dependent differences, and an exceptional set of mass $o(1)$. The last problem already has the flexible exponent by the discarded-remainder construction. Either a matching obstruction for the first problem or a way to rearrange the remainders in the second would resolve a precise limitation of the proof.

Research question V.54.2. What are the best leading constants in the flexible partition theorem for a fixed q and anisotropic tolerances?

The universal gap is less than four between the constructive and probabilistic constants. Geometry of numbers may improve the rectangular approximation step, while more careful distribution of chain remainders may improve the partition constant. The exact finite obstruction certificate provides a practical benchmark for small moduli.

V.54.2 Genuinely simultaneous polynomial recurrence

Research question V.54.3. For degree $k \geq 2$, which hypotheses permit simultaneous partition exponents polynomial in the number of phases, in place of repeatedly taking K th roots as in Lemma 5.9?

The linear proof cannot simply be reused by approximating all polynomial coefficients: after restricting to an arithmetic progression, the degree- j term contains d^j , and translation generates coupled lower-degree coefficients. A focused first case is a family of quadratics sharing a leading coefficient or lying in a bounded-height coefficient span. One should seek a joint recurrence statement tailored to those dependencies and then prove the same complete-partition interface. Success only for generic independent quadratics would still be valuable, but must be distinguished from a uniform all-coefficient theorem.

Research question V.54.4. Can a bounded-height generating family be chosen efficiently so that (I.26.6) is close to optimal for a given collection of modular phases?

The present theorem takes a representation matrix as input. Finding a better representation resembles an integer basis problem, with the analytic objective depending on both dimension and coefficient size. It is not ordinary vector-space rank. A useful output would be a certified generator matrix together with optimized positive budgets t_j , so every numerical gain is accompanied by a verifiable inequality.

V.54.3 Extremal density selection with additional information

Research question V.54.5. What is the exact largest guaranteed high-average *cell size* for a fixed number of cells, rather than the high-average mass solved by theorems I.27.1 and I.28.3?

Dividing mass by $M-1$ is useful, but it ignores how many cells are needed for negative mass and for the threshold cell in an extremizer. This is a finite constrained optimization problem. The natural first cases are $M=2, 3, 4$, followed by a piecewise formula in M . It may produce an improvement directly compatible with Lemma 5.15 without changing any analytic input.

[write] **6 October 2026; sources 11–14.** *Answered in the weighted model* by source 06's Theorem I.21.1: with M fixed and the cell masses free, the exact optimum is $\min\{P_+/(b(M-1)), S/(M-2)\}$ with explicit extremizers. Integer cardinalities and prescribed cell sizes remain open (item I.6).

Research question V.54.6. Can the joint frontier retain the entire list of phase radii s_j , cell weights, and a measured value of $V = \mathbb{E}|f|$, rather than a common radius and the endpoint upper envelope?

The affine aggregate inequality retains V exactly, while the sharp joint theorem optimizes over all functions in $[-a, b]$. Neither necessarily dominates the other after additional measurements are fixed. A dual optimization using a family of supporting lines to the distinct functions h_{s_j} is a concrete approach. An equality classification would show which measurements are actually useful.

Research question V.54.7. Which extremizers for theorem I.28.3 are ruled out when z is a polynomial phase on arithmetic progression cells?

The abstract extremizer places positive and negative function values at two phase endpoints. Algebraic phase constraints may prevent that configuration from occurring on many large cells simultaneously. Quantifying this obstruction would improve an application using genuinely new structural information, rather than attempting to strengthen an already sharp abstract inequality.

V.54.4 Cube geometry and higher-order cancellation

Research question V.54.8. Can the coloopless-subset polynomial $\sum_j c_{d,p}(j)z^j$ be evaluated or bounded effectively as d grows, with all characteristic dependencies explicit?

The quartic coefficient has the closed form (III.31.2), and dimension three is completely enumerated. Higher coefficients are a finite matroid problem. Complementation identifies coloopless sets with complements of flats of the dual matroid, suggesting a way to organize enumeration without examining every subset independently. Characteristic-dependent exceptional minors must remain visible. Even usable coefficient upper bounds could improve the crude binomial majorant.

Research question V.54.9. For fixed $\delta, \rho, \eta, \sigma$, what is the best possible upper value of $\mathcal{Q}_3(\delta + f)$, and which combinations of T_5, B_1, B_2, B_3, T_7 are jointly realizable?

The exact identity exposes the variables that were hidden by the norm triangle inequality. The separate estimates are not a full solution of their joint feasible region. The spectral deficit is attained by flat nonzero spectra, whereas the negative B_3 example uses a very different spectrum. Understanding that incompatibility is a plausible route to stronger bounds. A first target is to classify equality in $|B_2| \leq B_1$ simultaneously with $|T_7| \leq \eta^4 \sqrt{B_1}$.

Research question V.54.10. How far can the low-order terms in higher cube expansions be expressed using lower uniformity norms and a small list of spectral moments?

The three-cube calculation isolates an ordinary Fourier contribution at order four, a doubled-frequency moment at order five, and three six-point moments. For dimension four, one could first classify five- and six-circuits and ask which are controlled by U^2 rather than U^4 . This connects an explicit finite expansion with the theory of true complexity [29]; an effective inequality needs both the algebraic classification and quantitative analytic control.

V.54.5 A full quantitative and formal audit

Research question V.54.11. After replacing only the genuinely linear and real-valued local steps, what is the best complete bound obtainable by the existing Gowers proof architecture?

A valid answer requires a dependency-by-dependency parameter audit. Track the density at every iteration; the phase correlation before each partition; the number, dimension, and coefficient sizes of the phases; all exceptional masses; the actual new cell length; and the final transfer from modular to integer progressions. The improved telescoping count can relax a uniformity threshold, while the density-sensitive Fourier increment can improve local length. The higher-degree inverse step may nevertheless dominate the final asymptotic. Until that audit is complete, the correct output is a set of proved local improvements, as given here.

Research question V.54.12. Can the finite sharp statements and their extremizers be formalized as reusable Lean lemmas, with exact source normalization and an auditable numerical ledger?

This is a concrete implementation project. Begin with weighted real selection, then integer partition construction, then the normalized Fourier wrapper. For the cube identity, separate finite linear algebra from Fourier summation. Export the resulting constants and use the existing source ledger to distinguish completed companion statements from conditional consequences. The goal is a small set of stronger, reusable mathematical interfaces whose application to the larger proof is explicit.

V.55 Reproduction and source provenance

Source 13, delivered Appendix A.

The package contains the complete L^AT_EX source, compiled PDF, two standalone vector figures and their rendering script, four verification scripts, and their JSON outputs. Build the article from the package root using

```
python make_figures.py
pdflatex -interaction=nonstopmode -halt-on-error gowers_local_refinements.tex
pdflatex -interaction=nonstopmode -halt-on-error gowers_local_refinements.tex
```

The supplied figures are already rendered, so the first command is needed only to regenerate them. The density and partition scripts use the Python standard library. The cube and Fourier/counting scripts also use NumPy. The figure script uses NumPy and Matplotlib. The accompanying README gives the four verification commands. No original journal PDF or full third-party paper source is redistributed in the package.

The source snapshot was inspected through the user-selected GitHub connection. Its adapted paper has blob SHA 41085a0efde8c1932e86e80c791984f52841c22e; the status ledger has blob SHA feb6d45ba9e4a01092b8cf68083e1fbb99a1bf45. The accompanying provenance file records the pinned commit, relevant file paths, and hashes of the deliverable files. The article’s source comparisons refer to that snapshot. The mathematical proofs themselves are independent of future repository edits.

V.56 The Lean source map and a concrete formalization route

Source 14, delivered Section 8.

V.56.1 The inspected revision

The supplied repository [70] was inspected at commit

9ea383e9b28d885a936d5d681cd5584fffd2cb29.

The paper source and selected Lean files are recorded in the accompanying `source_manifest.json`, including upstream paths, byte sizes, and Git blob identifiers. The paper source matches blob 41085a0efde8c1932e86e80c791984f52841c22e. The source map below specifies where each refinement connects to the development.

Source location or declaration	Connection to this report
Definitions.lean, cubeCount, uniformity and cube-form definitions	Normalization bridge: divide the three-cube count by N^4 , use $f = 1_A - \delta$, and identify the uniformity budget with $\ f\ _{U^3}^8 \leq \alpha$.
Sections01_03.lean, lemma_3_10	The comparison point for Equation (F.19.1). A strengthened declaration must explicitly add the real centered and coprime-order assumptions, and its small-uniformity range.
Sections12_13.lean, DArrangement, arrangementCount	The two-coordinate arrangement model motivates the common-sidelength moment count. The present notation separates cube dimension from the number of cubes in an arrangement.
Sections14_15.lean, respectedCubePairCount, respectedGeneralArrangementCount	Direct counterparts of the collision count and the R_j family when all coordinate, index, and value groups specialize to $\mathbb{Z}_N$. The source's alternating-sign convention differs by a global sign, which preserves equality of label sums.
Sections14_15.lean, lemma_14_7	The $j = 8$ specialization of the moment step. Retaining $S = \sum_r B_r ^3$, and then $S - \Xi$ or S_{card} , provides a sharper local denominator.
Proofs09Moments.lean, Proofs12Parallelograms.lean, Proofs14Arrangements.lean	Existing moment and arrangement proof infrastructure to inspect when implementing the scalar interpolation and count identities. Our new stability and equality results are not already certified by merely referring to these files.
Sections08_09.lean, HasMonochromaticAP, corollary_8_3 [write: <i>HasMonochromaticAP is at Definitions.lean, line 113; corollary_8_3 is at Sections08_09.lean, line 52, and is open at HEAD e1db00e99</i>]	These encode monochromatic integer progressions. The polynomial supersaturation result requires a new finite-field reflection predicate and should not be presented as the same assertion.

Source location or declaration	Connection to this report
<code>Audit.lean</code>	The source audits theorem axiom dependencies against <code>propext</code> , <code>Classical.choice</code> , and <code>Quot.sound</code> , and counts exact companions for 120 numbered catalogue definitions. Its own comments distinguish this check from discharging conditional hypotheses or validating correspondence with the paper.

The repository also explicitly adjusts the printed exponent in Corollary 14.6 and propagates that adjustment to Lemma 14.8. The new moment denominator should be composed with the inspected statements and their hypotheses; it should not silently restore a printed exponent that the development has already corrected.

No Lean executable was available in the execution environment for this report. We inspected the source, but did not run the project build or the axiom audit, and did not add new Lean proofs. The article’s proof status is therefore conventional mathematics with reproducible finite diagnostics. This distinction is essential: a proposed theorem statement, an audited dependency boundary, and a complete kernel-checked proof are three different artifacts.

V.56.2 Recommended implementation order

The most direct first formalization target is the energy stability theorem. It requires finite sets, translations, counting, subgroup cosets, and rational inequalities, and does not depend on asymptotic notation or analytic Fourier transforms. A useful decomposition is:

1. Prove the overlap identity and the three-halves difference-set lemma for finite subsets of an additive commutative group.
2. Formalize the extraction of B and $H = B - B$, keeping $9/10$, $4/5$, and $5/4$ as rational constants.
3. Establish the two disjoint failed-triple counts in Equation (II.43.5). This is the essential combinatorial step and deserves its own reusable declaration.
4. Prove the polynomial defect inequality before introducing the square root. One can first state the conclusion as $d(1 - d) \leq \varepsilon$ with $0 \leq d < 1/3$, then derive the τ formulation separately.
5. Add the coordinate-corner homomorphism argument. Its image and fiber cardinalities force $e_i H \subseteq H$ without a dimension-dependent estimate.
6. Formalize the exact labeling classification using finite differences and inclusion–exclusion anchored at the origin.

The second target is the finite moment equality theorem. The scalar convexity argument can be separated from Fourier orthogonality. After obtaining the zero-or-common-modulus Fourier criterion, nonnegativity and equality in the triangle inequality recover constant weights on cosets. The combined geometric classification then follows through the zero sidelength and graph supports.

The cube expansion is naturally divided into a finite support classification, the Fourier identity, a finite inequality for S, D, ρ, b , and a real-variable optimization lemma. The asymptotic theorem

should be a final wrapper around these exact statements. This avoids encoding a long argument entirely in little-oh notation.

For polynomial colorings, begin with the reflection-coefficient equivalence and its Vandermonde hypothesis $p > k$. The minimum-support lemma, equality slicing, and norm construction are reusable finite-field results in their own right. Definitions should keep the palette $\mathbb{F}_q^r$, its coordinate degree bounds, ordered parametrizations, and the zero direction explicit.

V.57 A focused research program

Source 14, delivered Section 9.

The following priorities identify concrete next results and the information each would add. These are proposed research questions; they are not asserted to be established open problems with an exhaustive priority history.

1. **The optimal exponent-five coefficient.** For nontrivial exponent-five groups, determine whether a universal coefficient C_5 gives

$$M_{\delta,G}(u) = \delta^8 + 6\sqrt{2}\delta^4 u^4 + C_5 \delta^2 u^6 + o_\delta(u^6),$$

and whether $C_5 = 2^{3/4}/3$. The proof supplies the sharp exponent and a candidate model, but arbitrary small Fourier tails and their phases must be controlled before identifying the coefficient. Dependence on the rank of G is itself part of the question.

[write] **6 October 2026; sources 21–27.** *Answered by sources 21 and 23:* the universal coefficient exists and $C_5 = 2^{3/4}/3$, with an absolute $O(u^8)$ remainder independent of the rank (source 21's Theorem F.26.1, source 23's Theorem III.58.1); it is attained on every fixed group (source 21's Corollary III.55.2). Item III.23.

2. **The next term outside exponent five.** Determine the first term after $C_* \delta^{8/3} u^{16/3}$. The order-seven computation already warns that the next answer can depend on torsion. A plausible initial project is to compare fixed groups with minimum nontrivial character order at least eleven against order-seven groups, retaining every Fourier resonance at the first unresolved scale.
3. **Effective secondary rigidity.** Replace the sequential conclusion of Theorem III.34.3 by an explicit estimate in terms of the normalized deficit and u/δ . The final scalar optimization has a nondegenerate maximum. This should convert the proof into quantitative control of amplitude, phase, and unused Fourier fourth mass. A separate question is what additional hypothesis permits an L^2 approximation.
4. **The largest range of the energy envelope.** Determine the best universal range of ε for $|A\Delta C| \leq \tau(\varepsilon)|A|$, and separately the range forcing the nearest coset to split by coordinates. The current $1/100$ is a safe explicit range of the extraction argument. Classifying every exact envelope example would sharpen the present infinite-family sharpness statement.
5. **Quantitative repair of labels.** Once the domain is close to a product coset, turn the collision defect into a Hamming-distance estimate for $M + \sum_{U \subseteq [k]} f_U$. Dependence on the number of coordinates, the target group, and degenerate coordinate factors should be stated explicitly. A useful first case is a full $\mathbb{F}_p^2$ domain with labels in $\mathbb{F}_p$.

6. **Stability of simultaneous arrangement equality.** Find hypotheses under which near equality in the profile-moment bound yields a common coordinate subgroup, an approximately coset-supported family of section indices, and approximately $(k + 1)$ -additive top interactions. The spectral variance estimate controls weighted Fourier mass, but does not alone provide this joint spatial description.
7. **Dimension thresholds above the reflection budget.** Resolve the scalar degree-five, symmetric four-term problem, beginning with dimension five. The present bounds are $4 \leq d_{\max} \leq 8$ for the specified model and characteristic range. The Artin–Schreier example rules out one tempting norm construction in characteristic five; it does not rule out other quintic maps.
8. **Classification of exact polynomial extremizers.** Describe all cubic maps at the threshold in dimension four and all possible center-dependent subspaces $S_z(Q)$ above it. The equality geometry of the zero set is settled; compatibility as the center varies is the missing global condition. The example $x + y^2$ excludes a universal fixed-kernel answer.
9. **Second-minimum direction counts.** Determine the next attainable value of $|S_z(Q)|$ in dimensions $d \geq B + 2$, and over nonprime fields in dimension $B + 1$. The zero sets come from homogeneous odd coefficients of a single polynomial map, which may impose more restrictions than an arbitrary bounded-degree system.
10. **Palette size and redundant coordinates.** Replace the prescribed degree-vector budget by a useful invariant of the polynomial map’s actual image. The current theorem counts coordinates with degree bounds, even when some color information is redundant. Any improvement should separate algebraic redundancy from arbitrary relabeling of the palette.
11. **Higher-dimensional centered cube extremals.** Identify the first surviving centered correlations for U^d with $d \geq 4$, determine the sharp real mean-zero comparison with lower uniformity norms where one exists, and study the torsion-dependent extremizers. A successful method must control how different support types compete, not merely bound each one separately.
12. **A complete parameter ledger for the global argument.** Insert the profile and geometric denominators into an entire arrangement-extraction and density-increment chain. Track the densities, multiplicities, degeneracies, and threshold choices symbolically. The relevant output is a proved end-to-end parameter improvement; a better isolated local constant is not enough to infer it.

The energy stability theorem and the finite equality classifications are suitable near-term formalization projects. The order-five coefficient and the quintic coloring problem are compact research problems with explicit candidate constructions and obstructions. The higher-cube and global-density questions require additional ideas beyond the local proofs developed here.

V.58 Reproducibility, diagnostics, and proof status

Source 14, delivered Appendix B.

V.58.1 Contents of the source package

The package contains one complete article source, `gowers_extremal_refinements.tex`, its compiled PDF, the PDF figure and figure-generation script, two verification scripts and their recorded JSON results, the pinned source manifest, and a short build guide. The external paper and the Lean repository are referenced, rather than redistributed. The bibliography is embedded in the article source, so no BibTeX run or separate bibliography database is needed.

The article compiles with a standard pdfLaTeX installation:

```
pdflatex -halt-on-error gowers_extremal_refinements.tex
pdflatex -halt-on-error gowers_extremal_refinements.tex
```

The existing figure is included, so Python is not needed to build the PDF. To regenerate the computational records and the figure, use:

```
python3 verify_results.py
python3 verify_certificates.py
python3 make_figures.py
```

The first script requires NumPy, the certificate script uses only the standard library, and the figure script requires Matplotlib. All files are addressed relative to the scripts.

V.58.2 Checks performed

Check	Scope and result
Sparse Fourier identity	All $5^8 = 390\,625$ assignments of frequencies in $\{-2, -1, 0, 1, 2\}$ were tested against the four integer orthogonality constraints. Exactly 225 survive, in the eight histogram rows of the model table.
Centered support certificate	All 255 nonempty subsets of the cube vertices were classified with exact rational arithmetic. The quartic and quintic surviving counts are 12 and 8.
Small-torsion certificates	Every Fourier slot in groups of orders 5 and 7 was checked as an exact polynomial, including the phase cancellations in order five. All table entries and cube polynomials match.
Fourier inequality diagnostics	144 real spectra were tested on cyclic groups of orders 5, 7, 11, 13, 17, 25, including concentrated spectra and an order-five dominant character inside a larger group. The direct cube and Fourier formulas agree with maximum normalized numerical discrepancy below $2.5 \cdot 10^{-16}$. All tested norm inequalities pass the stated tolerances.
Energy and algebra	All 2036 nonempty subsets of cyclic groups of orders 1 through 10 were checked for the energy upper bound and coset equality. The defect-polynomial identity and nonnegativity were checked in 35 300 rational cases. Five larger prime examples verify exact sharpness.

Check	Scope and result
Coordinate geometry and labels	All 511 nonempty subsets of $\mathbb{F}_3^2$ were checked: exactly 16 attain the cube-collision maximum, precisely the coordinate cosets. Among all 19 683 labels of the full domain into $\mathbb{F}_3$, exactly 729 attain equality, matching $axy + u(x) + v(y)$.
Polynomial norm coloring	The norm associated with the irreducible polynomial $T^3 + T + 1$ over $\mathbb{F}_5$ was checked on all 125 inputs. For its cubic four-coordinate coloring, all 625 centers and all 625 directions per center were checked. Only the zero direction survives at every center.
Asymptotic arithmetic	The exact two-harmonic norm and cube formulas were evaluated at $u = 10^{-1}, \dots, 10^{-9}$ with 90-digit decimal arithmetic, verifying the norm constraint and convergence of the secondary residual to C_* .

The randomized diagnostic seed is 20261006. Numerical tolerances are recorded in `verification_results.json`; the exact tables are in `certificate_results.json`. Finite enumeration can detect normalization errors and missed torsion cases, but it does not establish the all-group inequalities, uniqueness statements, or asymptotic uniformity. Those claims are established by the proofs in the main text.

The principal arguments were also reviewed independently within the preparation process. The review led to explicit handling of order-five coincidences, the phase-error term in secondary rigidity, the small-field boundaries in the minimum-support proof, and the distinction between geometric support repair and label repair. This is internal review, not external peer review or formal verification.

V.59 Reproducible computations and Lean integration

Source 15, delivered Section 13.

V.59.1 Finite checks supplied with the manuscript

The accompanying `verify.py` uses only the Python standard library and a fixed random seed, 20261006. It computes arrangement counts by convolution, not by enumerating all sixteen cubes of an eight-arrangement. For each h , let m_h^{*d} be the d -fold additive convolution on $T \times V$. The exact identity

$$R_d(h) = \sum_{z \in T \times V} (m_h^{*d}(z))^2 \quad (\text{V.59.1})$$

is simply the count of two ordered d -tuples with the same sum. For the tested integer and rational measures, this uses exact arithmetic.

The executed checks include all 1,036 nonempty subsets of $(\mathbb{Z}/2\mathbb{Z})^2$, $(\mathbb{Z}/3\mathbb{Z})^2$, $(\mathbb{Z}/2\mathbb{Z})^3$, and $\mathbb{Z}/4\mathbb{Z} \times \mathbb{Z}/2\mathbb{Z}$. They verify the cubic cap and the equivalence between its equality cases and Cartesian cosets. Most stability hypotheses in these very small groups occur only at exact equality. Two additional nonzero-deficit examples, in $\mathbb{Z}/101\mathbb{Z}$ and $\mathbb{Z}/101\mathbb{Z} \times \mathbb{Z}/2\mathbb{Z}$, have deficit 99/10000 and produce the asserted coset certificates.

The joint normal form is checked against the count equalities for all 6,560 nonempty partial binary labellings on $(\mathbb{Z}/2\mathbb{Z})^2 \times \mathbb{Z}/2\mathbb{Z}$, and for 180 further ternary cases. Other tests cover random indicator domains, rational vertex weights, sharp subgroup examples, exact label decompositions, the polynomial stability identity, and all displayed exponent comparisons over a finite range. The weak parameter $q^{-1/8}$ in Lemma II.49.2 supplies a valid product property for the nonconstant random labels; the coupled inequalities are checked with its powers represented exactly as rational numbers.

Thirteen Fourier cross-checks use complex floating-point arithmetic with relative tolerance 10^{-9} . They are deliberately distinguished from the integer and rational count tests. The exact counts and all inequalities are not certified by those floating-point computations. The complete executed inventory is in `verification_results.json`, with a readable log in `verification.txt`. Every assertion passed in the supplied run. Finite examples supplement the proofs and do not establish a universal statement by themselves.

V.59.2 An explicit constructive certificate for the stability theorem

Given the finite set A , compute $r(t)$ and $P = \{t : 8r(t) > 7|A|\}$. Form $H = P - P$ and choose a coset of H maximizing its intersection with A . Under the theorem’s hypothesis this construction returns the asserted Cartesian coset. A verifier can check subgroup closure, coordinate-projection invariance, the two cardinality bounds in (II.52.1), and the final symmetric difference directly. The supplied script implements this construction. After the correlations are available, subgroup and coset construction require at most quadratic group-table work, with a separate coordinate-invariance check. Computing the input cube count naively has a 2^k vertex cost.

V.59.3 Repository interfaces that would support the new results

The inspected proof of Lemma 14.7 already has a general even-moment identity and the interpolation infrastructure, but the relevant helpers are private [80]. The shortest development path is therefore not to [6 *October 2026; sources 15–20; checked at HEAD 0b4890a1c*] *Precision*: the general even-moment identity is the private `cor147_moment_eq_count` (`Proofs14HigherArrangements.lean`, line 484, general d); the interpolation helpers `cor147_interpolation` and `cor147_moment_interpolation` (lines 653, 701) are hard-wired to $d = 8$. The unweakened Corollary 14.6 exponent is already kernel-checked as the private `cor146_direct_lower` (note in Section F.20).

formalize an unrelated Fourier theory anew. It is to expose an appropriate public interface and retain the intermediate count.

A practical module sequence is as follows.

1. **Public moment identity.** Export a general- d version of the moment–arrangement correspondence. Define the collision parameter using `respectedGeneralArrangementCount 1 B phi`, avoiding an unnecessary second notion of the same count.
2. **Division-free interpolation.** Prove $R_2^{d-1} \leq R_1^{d-2} R_d$ over real casts. Treat $R_1 = 0$ before introducing normalized quotients. The case $d = 2$ should be separated from strict-convexity arguments.
3. **Fibre cap.** Formalize the injection $(h, y, z) \mapsto (y, y + h, z)$ and then the label-forgetting injection. Only finite group arithmetic and cardinalities are required.

4. **Collision-parametrized product step.** Refactor the reference-cube argument in `Proofs14Arrangements.lean` to conclude $R_2 \geq \gamma^{u_k} C^4 / (nM^7)$ before inserting the density lower bound. Its documentation already makes clear that the reference base has to be retained in an injection or accounted for by averaging [81].
5. **Coupled exponent arithmetic.** Combine the previous two public results and only then substitute the collision lower bound. Introduce a new stronger theorem and recover the existing corrected statement as a corollary, rather than silently changing the meaning of an existing name.
6. **Independent stability modules.** The small-difference lemma, energy-deficit argument, coordinate-mixing argument, and exact joint normal form can be developed separately from the analytic proof chain. They should have their own statements and checks rather than being counted as completed source lemmas.

The source checkpoint reports a successful build and axiom audit of its existing development [83]. That report was read, not reproduced here. This package contains no newly kernel-checked Lean theorem and makes no claim to close later open stages in the repository. The accompanying `integration.md` records proposed interfaces and the same limitations. In particular, a mathematically proved enhancement is not already a Lean enhancement merely because its intended dependencies exist.

V.60 Relation to existing work and remaining research

Source 15, delivered Section 14.

V.60.1 Literature position

The basic moment conversion is Gowers’s mechanism, and the near-maximal energy part of the Cartesian proof belongs to the familiar almost-period approach in additive combinatorics. Our proof adds an explicit coordinate mixing argument because an arbitrary subgroup need not split into coordinate subgroups, as Example II.53.4 demonstrates.

Near-extremizers for ordinary Gowers–Host–Kra seminorms are studied by Eisner and Tao [7]. Kovač and Rogers’s September 2026 preprint studies sharp bounds, near-extremizers, and degree-lowering inequalities for L^p norms of multiple autocorrelations [71]. These are relevant surrounding theories, but the present functional counts *coordinate-restricted axis cubes with a shared side vector*; it is not being identified with the usual unrestricted Gowers norm. In particular, results whose hypotheses exclude compact open subgroups cannot simply be applied to finite groups to improve a constant that is attained by the subgroup examples above. This manuscript claims neither to supersede that literature nor to settle its already resolved questions.

The investigations below are concrete follow-on questions generated by the proved results. They are not represented as a list of historically recognized open problems. Some may admit a reformulation using existing stability or property-testing results.

Question V.60.1 (Best Cartesian stability constant). *What is the smallest uniform function $\delta(\varepsilon)$ for which $Q(A) \geq (1 - \varepsilon)|A|^3$ forces a Cartesian coset within $\delta(\varepsilon)|A|$ points?*

Theorem II.52.2 gives $\delta(\varepsilon) \leq 18\varepsilon$ up to $1/100$, while Example II.52.3 gives $\liminf_{\varepsilon \rightarrow 0} \delta(\varepsilon)/\varepsilon \geq 1$ along an admissible sequence. Determining the leading constant and the largest valid near-extremal range is a finite, explicit extremal problem.

[write] **6 October 2026; sources 15–20.** *Leading constant answered by source 14: its Theorem II.43.3 gives $\delta(\varepsilon) \leq \tau(\varepsilon) = \varepsilon + O(\varepsilon^2)$ up to $1/100$, and Example II.52.3 (= source 14's Proposition II.43.4) attains it, so the leading constant is 1, not 18. The largest valid range is source 14's Question 4 and stays open (item II.31).*

Question V.60.2 (Quantitative repair of labels). *Suppose the respected cube-pair count is at least $(1 - \varepsilon)|A|^3$. After passing to the nearby Cartesian coset, can the labels be altered on $O_k(\varepsilon)|A|$ points, with appropriate treatment of missing points, to obtain the exact form (II.51.1)? Can the dependence on k be removed?*

The allowed correction class must include all coordinate-omitting functions and must work in arbitrary finite abelian value groups. Requiring only an ordinary polynomial would miss the exact examples already classified. The denominator measuring errors should be specified before attempting a property-testing argument on sparse fibres.

Question V.60.3 (Robust joint normal forms). *If both $1 - C/\Gamma$ and $R_8 C^6/R_2^7 - 1$ are small, is there a quantitative analogue of Theorem II.54.2 in which most fibres share one horizontal subgroup and the top mixed values are close to a $(k + 1)$ -multiadditive model?*

The spectral theorem supplies weighted closeness to a common magnitude; the geometric theorem supplies separate fibre cosets. Their combination still requires coherence across fibres. A robust conclusion must retain arbitrary fibre translates and coordinate-omitting terms, and must state whether the approximation is in point mass, collision mass, or cubic mass.

Question V.60.4 (Preserving the coupled gain in the later proof). *How much of the improvement in (II.50.4) survives the selection and density-increase arguments in Sections 15–17, and the quantitative iteration in Section 18?*

A useful first deliverable is a symbolic parameter ledger carrying $(\beta, \gamma, \chi, \rho_3)$ through each extraction. Replacing these by single worst-case density parameters too early could erase the gain. Only a complete ledger with verified hypotheses could justify a statement about a final Szemerédi bound. The present article does not provide that last step.

Question V.60.5 (Optimal product-property amplification). *Are the powers $3d - 2$ of χ and the associated exponents in (II.49.7) optimal under PP_γ alone, or can the reference-cube averaging exploit further compatibility between side vectors?*

The constants in raw moment amplification are sharp, but its simultaneous combination with the configuration lower bound may have incompatible extremizers. Establishing a genuinely stronger theorem would require examining that compatibility, not just optimizing an isolated application of Jensen. The side-vector profile in Proposition II.53.3 is an explicit place to begin.

Question V.60.6 (Balanced extraction versus density cost). *When can one pass to a substantial subdomain with nearly equal fibre densities while retaining enough respected arrangements to benefit from ρ_3^{-6} or the precise profile bounds?*

An answer should optimize the retained arrangement mass against the loss in β and the number of occupied fibres. Example II.48.4 shows why a balance assumption cannot simply be inserted without a cost.

Question V.60.7 (Weighted near-extremizers). *Can near equality in the weighted cubic cap be converted into a useful rounding theorem for f , under natural lower-density or nonconcentration hypotheses?*

The exact argument in Section II.56 proves only that exact saturation forces zero-one weights. Controlling the total mass of fractional weights from a global cubic deficit is a separate issue. Its resolution would connect the weighted coupling theorem to geometric stability without imposing an indicator-domain reduction at the outset.

Question V.60.8 (Nondegenerate and anisotropic variants). *What error term is necessary when arrangements are required to have distinct vertices, or when the allowed side vectors come from a prescribed anisotropic family?*

A fixed side restriction preserves the moment identity, whereas excluding cross-cube coincidences need not. A satisfactory theorem should distinguish these cases and quantify the exceptional counts, especially when the ambient groups have small torsion. Prime cyclic asymptotics should not be used to justify a statement for arbitrary finite groups.

Question V.60.9 (Certificates and formal interfaces). *Can the public moment identity, collision cap, and reference-cube bound be assembled into short checked Lean proofs of the stronger statements, with an automated exponent ledger and executable finite certificates?*

The immediate target is the division-free all- d inequality followed by the collision-parametrized product step. The structural stability results can then be checked independently. This separation would make the distinction between analytic improvements, combinatorial models, and final iteration obligations explicit in the proof status.

[write] **6 October 2026; sources 15–20.** *Cross-answers.* The second and third questions above are source 14’s Questions 5 and 6 (item II.24); the fourth restates source 14’s Question 12 and source 04’s Question 7 (items V.1, II.3), and the sixth is a variant of source 04’s Question 7. The others are new in this report; all are merged in Section V.2.

V.60.2 A proposed order of work

The highest-leverage next step is the collision-parametrized Lean interface, because its proof ingredients are already present and it records the main quantitative saving. In parallel, the elementary Cartesian stability proof is a contained formalization project. A parameter ledger through later sections and a robust version of the joint normal form are the more substantial research tasks. Neither should be replaced by a claim that the local gains automatically constitute a global breakthrough.

V.61 Conclusion

Source 15, delivered Section 15.

The arrangement argument loses information at two different places: in replacing the collision count by an ambient second-moment cap, and in eliminating that collision parameter before combining the product-property step with interpolation. Keeping it produces sharp density-sensitive amplification and the stronger coupled bound

$$a_8 \geq \beta^{22 \cdot 4^k} \gamma^{44k4^k + 56 \cdot 2^k}.$$

The same parameter also has a geometric meaning. Its cubic cap has exact Cartesian extremizers, its near-extremizers satisfy a dimension-free linear-error theorem, and simultaneous moment and cubic equality have the complete fibrewise multiadditive normal form in Theorem II.54.2. These conclusions are proved independently of the uncompleted later stages of the supplied formalization.

V.62 Notation and verification inventory

Source 15, delivered Appendix A.

Symbol	Meaning
X, M	Horizontal product group and its cardinality.
T, n	Vertical group and its cardinality.
V, q	Value group and its cardinality.
B_x, b_x, β_x	Horizontal fibre, its size (or mass), and its normalized density.
β, ρ_3	Mean fibre density and mean cubed fibre density.
$R_d, R_d(h)$	Respected d -arrangements, globally and at a fixed side vector.
C, C_h, C_x	R_1 and its side-vector and vertical-fibre decompositions.
Γ	Cubic cap $\sum_x b_x^3$.
a_d, χ	Counts normalized by $M^{2d+1}n^{2d-1}$ and M^3n .
$Q(A)$	Unlabelled parallel axis-cube-pair count in A .
e_k, u_k	$2k4^k$ and $8 \cdot 2^k$, the two product-property exponents.
$\mathcal{D}_d$	Relative defect $R_d C^{d-2} / R_2^{d-1} - 1$.

Table V.62.2: Executed finite checks. Counts refer to cases, not to individual assertions within a case.

Family	Cases
Exhaustive nonempty horizontal subsets	1,036
Additional positive-deficit stability certificates	2
Random indicator arrangement domains	48
Constant-label coupling tests	48
Rational vertex-weight domains	12
Sharp subgroup moment model	1
Exhaustive binary label maps on $\mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/3\mathbb{Z}$	64
Additional random label maps	48
Multiadditive composite-modulus label models	4
Exhaustive partial-domain joint-extremizer cases	6,560
Additional ternary joint-extremizer cases	180
Exact rational spectral-stability cases	320
Integer exponent comparisons	84
Floating-point Fourier identity cross-checks	13

The script also checks the unbalanced-fibre counterexample and the diagonal subgroup counterexample. All random choices are deterministic from the recorded seed. No external mathematical package or network service is required to reproduce the checks.

V.63 Source pin and proof-status boundaries

Source 15, delivered Appendix B.

The repository was inspected at the immutable commit

9ea383e9b28d885a936d5d681cd5584fffd2cb29.

The supplied paper TeX has blob identifier

41085a0efde8c1932e86e80c791984f52841c22e.

The consulted original PDF has printed Lemma 14.7 and Lemma 14.8 on page 560; the arrangement definition appears on page 559. The repository’s Section 14 definition file was inspected separately because its repaired exponents differ from the printed statements. The exact URLs, blob identifiers available for the inspected files, and access date are recorded in `provenance.json`.

The package contains a mathematical manuscript, compiled PDF, executable finite verification, and an integration plan. It does not contain a new compiled Lean proof, a patch pushed to the repository, or a proof of any new global bound for progression-free sets. Its universal conclusions rest on the written proofs, not on the finite tests or on an unperformed build.

V.64 How the refinements enter the larger argument

Source 16, delivered Section 9.

The three improvements affect different interfaces, and their quantitative information should be kept separate until an actual composition is justified.

V.64.1 From a local phase to a smaller cyclic model

For a cell Q of upper size U , a lower bound

$$\mathcal{C}_{k+1}(\mathbf{1}_Q g) \geq \beta U^{k+2}$$

translates, under a cube-preserving identification with a subset of a new cyclic group of order p' , into

$$\|g_Q\|_{U^{k+1}}^{2^{k+1}} \geq \beta (U/p')^{k+2}. \quad (9.1)$$

This formula is only a normalization identity once the identification preserves all cubes being counted. The existence of a suitable identification and the avoidance of wraparound must be proved for the chosen interval model.

Theorem F.21.1 supplies larger cells than the short-cell implementation in the source proof. Their maximum remains a controlled multiple $3km$, so the scale change is explicit. Whether a later step benefits from the improved factor requires substitution of this upper size into (9.1) and into all polynomial partition thresholds. The active-pattern theorem, Proposition IV.14.2, provides an alternative when exact preservation of the earlier scale is more important.

V.64.2 From affine spectra to local linearity

The moving-frequency condition $|(a_i h + b_i) d_Q|_N \leq \delta N$ places the actual progression step in the corresponding Bohr constraints. The certificate theorem constructs that condition directly. It can therefore be inserted at an interface which requires small frequencies times the step, without first converting a quadratic diameter statement.

The uniform recurrence threshold in (6.5) must still be discharged after the number q of affine pieces has been chosen. In the source, q can itself be a very large function of the density parameter. An improvement from exponential to quadratic dependence on q in the local exponent is substantial, but it is not a calculation of the complete final bound.

V.64.3 From exact cube coefficients to inverse information

The low-order cube formulas identify which small vertex configurations create the first terms of a balanced perturbation. They allow precise bookkeeping of low-dimensional Fourier structure inside a higher cube. They also show where torsion enters: order two already affects the fourth term, and order three adds new fifth-order configurations from dimension four onward.

Turning this information into an improved inverse theorem requires control of the remainder and a mechanism extracting useful structure from the relevant Fourier expressions. The coefficient identities alone do not provide that mechanism. They do provide exact targets against which a proposed inequality or a formal expansion can be checked.

V.65 Source fidelity and a formalization plan

Source 16, delivered Section 10.

V.65.1 What was actually inspected

The repository snapshot identified in Section F.21 reports 98 exact companion proofs and 22 open catalogue statements, with a separate axiom audit [84]. *[6 October 2026; sources 15–20] Stale since the pin:* at HEAD 0b4890a1c the ledger records 101 exact companions and 19 open statements (lemma_16_9, lemma_13_8 and lemma_13_7 have closed since dca9f0638), and FORMALIZATION_ST ATUS.txt reports 233 modules and 2,138 audited public theorems (169 and 1,658 at the pin). Those are the repository’s recorded checkpoint results; this report does not claim to have rerun that Lean build. The status file also distinguishes repaired signatures and conditional implications from full source-fidelity claims.

The following files are particularly relevant to the new work:

- `Sections17_18.lean` gives the explicit partition uniformity normalization and the statement of Proposition 17.7.
- `Proofs17Phase.lean` contains the polarization and phase-twisting infrastructure.
- `Proofs17LocalizedPhaseRemoval.lean` contains the existing carry-pattern implementation of the local step. *[6 October 2026; sources 15–20]* It is the proof of `proposition_17_7_holds` (line 2900), an exact companion at the pin and at HEAD (note in Section F.21); its small

branch is the private `prop177_small_branch` (line 1283), the argument of this source’s small case. The polarization (3.2) is the private `powersetDifference_pow` and the phase (3.1) the private `polarizationPhase` (`Proofs17Phase.lean`, lines 142 and 217), used by `lemma_17_1_holds` (line 374).

- `Proofs13QuadraticRecurrence.lean` isolates the Section 13 transport and weighted selection interface, including documented scale and endpoint issues.

These source observations motivated the exact $l + 1$ normalization, the properness conditions, and the explicit separation of finite certificates from external existence thresholds.

V.65.2 Suggested theorem boundaries

A formalization can be divided into reusable finite statements.

1. **Cyclic window geometry.** Construct the sets S_t in the enlarged cyclic group, prove their two possible cardinalities, prove the exact coverage count (4.4), and prove that each family of equally spaced centers gives a partition. No Fourier analysis is needed here.
2. **Phase-independent localization.** State Theorem IV.13.2 for arbitrary functions f_ϵ and an explicitly supplied polarization identity. This avoids coupling the geometric theorem to a factorial argument.
3. **Polynomial synthesis.** Prove (3.1) and the formal coefficient criterion. Over $\mathbb{Z}/N\mathbb{Z}$, keep factorial divisibility as a named hypothesis instead of hiding it in simplification.
4. **The exact certificate algorithm.** Formalize Theorem I.30.2 from a supplied s . The required operations are residue partitions, interval chunks, rectangular pigeonhole, and a triangle inequality in $\mathbb{R}/\mathbb{Z}$. The recurrence theorem can then be attached as a separate imported interface.
5. **Small cube configurations.** Represent cube vertices as Boolean functions, classify four- and five-element supports, and combine finite Fourier orthogonality with the integer relation modules described in Section III.39.

Kernel checking would certify these local implications. It would not automatically certify that a new final Szemerédi exponent follows from the existing higher-dimensional structural chain. The latter needs a separate theorem with all numerical hypotheses and model changes visible in its type.

V.66 Reproducible computations and their role

Source 16, delivered Section 11.

The archive contains three verification programs and their recorded output. They use small cases to challenge the algebra and geometry; the proofs in the preceding sections establish the general results.

Program	What is checked
<code>verify_localization.py</code>	Finite cyclic support containment, cyclic partitions, phase signs, exact interval cube counts, and direct localization inequalities.
<code>verify_partitions.py</code>	An exact rational recurrence certificate, full coverage by the constructed cells, cell lengths, and every resulting step inequality.
<code>verify_cube_coefficients.py</code>	Enumeration of small binary supports and exact rational cube-polynomial coefficients in several cyclic groups.

The localization energy calculation uses floating-point complex arithmetic, with recorded numerical tolerances; its combinatorial support checks use integer arithmetic. The partition and cube coefficient checks use exact arithmetic in their stated finite examples. None of these programs is a substitute for a proof of the cited recurrence theorem, a Lean check, or an exhaustive literature search.

V.67 Further research questions and priorities

Source 16, delivered Section 12.

V.67.1 Local phase removal

1. **Determine the best retention factor at controlled cell size.** Fix $C \geq 3$, an upper size Ckm , and a lower size $m/(3k)$, and ask for the largest coefficient $c_k(C)$ for which a statement of the form Theorem F.21.1 holds. The sliding construction and the local-scale ceiling determine the leading $k \log k$ loss. Close the remaining factor exponential in k , or find a sharper obstruction.
2. **Retain cells of size $O(m)$.** Can one obtain a factor $2^{-O(k \log k)}$ while also forcing every cell to have upper size Cm for an absolute C ? The active-pattern theorem preserves the short-cell scale, but still contains a global pattern count.
3. **Optimize nonuniform windows.** The present partitions use equal real widths. Input increments with unequal side lengths suggest anisotropic displacement budgets, and input energy concentrated away from extreme increments may permit smaller windows. A useful theorem would retain all weights through the localization step.
4. **Extend the polarization interface.** Theorems IV.12.2 and IV.12.4 settle the formal coefficient test and the classical prime-field functional test. Investigate corresponding localization theorems for nonclassical torsion-valued phases, stating their target groups and degree conventions explicitly.

V.67.2 Partitions controlling their own steps

5. **Close the dimension gap.** For equal powers, the sufficient exponent is of order q^{-2} , whereas Theorem I.31.1 permits at most order q^{-1} . Can a universal partition theorem attain c/q ?

6. **Determine the anisotropic feasible region.** Is the strict interior of $(q + 1)a + 2 \sum_i b_i \leq 1$ sufficient, perhaps with logarithmic losses? The present work proves necessity and a smaller sufficient region, not equivalence.
7. **Use the weight during construction.** A weighted extraction of one useful progression might beat the full-partition obstruction. Develop a direct weighted theorem suited to the arrangement counts in Lemma 13.5.
8. **Measure the effective spectrum dimension.** Find structural hypotheses ensuring that the moving slopes lie in a low-dimensional bounded-height integer span. Combine such a result with Corollary I.30.8.
9. **Make thresholds numerical.** Track the constants in a simultaneous recurrence theorem well enough to replace $X_0(C, q)$ by an explicit computable bound. Alternatively, formalize a certificate-producing finite algorithm and compare its actual performance with the worst-case existence theorem.

V.67.3 Cubes and the full quantitative induction

10. **Classify sixth-order supports.** Determine the relevant primitive relation modules for six Boolean vertices in arbitrary dimension. The fourth- and fifth-order calculations suggest a systematic classification by small integer matrices and their torsion quotients.
11. **Control the higher-dimensional remainder sharply.** The remainder estimates supplied by repeated Gowers–Cauchy–Schwarz are uniform but rarely optimal. Determine sharp dependence on lower Gowers norms, the density, and the torsion profile.
12. **Find extremizers and stability statements.** The coefficient formulas identify equality mechanisms at low order. Which functions nearly maximize the complete higher-dimensional cube polynomial under fixed lower-order norm constraints?
13. **Propagate the local gains through the complete proof.** Substitute the new partition scale and recurrence thresholds into the structural, localization, remodeling, and density-increment steps. Prove the resulting numerical recurrence before claiming an improvement to the final Szemerédi bound.
14. **Build a kernel-checked local refinement package.** Start with the finite certificate theorem and sliding-window geometry. They have explicit finite statements and avoid the currently open global dependencies. Then add the exact cube coefficients and a separate source-fidelity audit.

V.68 Conventions for reproducibility and citation

Source 16, delivered Appendix A.

The article is self-contained at the level of its new finite proofs. The mixed cube inequality is the standard Gowers–Cauchy–Schwarz inequality; the modern simultaneous recurrence theorem is cited as an external result. The supplied code records finite checks and can be rerun independently.

The files in the archive are:

- this article in L^AT_EX and PDF;
- the three verification programs and their output;
- the plot source, its formula-based comparison table, and the figure;
- a README with build and verification commands.

The source repository is referenced at a fixed commit so that a later change to its formalization status or statement repairs does not silently change the meaning of this report.

V.69 Relationship with the inspected Lean development

Source 17, delivered Section 9.

V.69.1 Snapshot and relevant interfaces

The repository inspection was anchored to commit

dca9f06383ccdfc2ed208ccb7c9882214fdb09ac.

The relevant directory is

Combinatorics/Ramsey/Lean/GowersSzemerédi.

The statements in `Sections08_09.lean` use `phiAdditiveTupleCount`, `additiveTupleCount`, `GammaAdditive`, and `GammaHomOfOrder` [85]. The inspected `Proofs09Moments.lean` proves the graph Fourier-moment identity and then Lemma 9.2 [86]. The restriction module imports `Proofs15Restriction` and presents Lemma 9.3 through the height-zero arrangement interface [87]. These are source observations, not claims that the entire repository was rebuilt during preparation of this note.

V.69.2 A particularly small first change

The existing proof of `section9_second_moment_1e` first obtains

$$\sum_{u,r} |\hat{f}_u(r)|^2 = N^2 |B|$$

and only then uses $|B| \leq N$. [6 October 2026; sources 15–20] `section9_second_moment_1e` is a private lemma (`Proofs09Moments.lean`, line 480); the equality $N^2 |B|$ is a step inside its proof, and the fourth and sixteenth moment identities exist only inside `lemma_9_2_holds` (line 527). The first formalization task is to expose the equality as its own reusable lemma. Combining it with the existing fourth- and sixteenth-moment identities should give the division-free target

$$G_{16}(B, \phi) |B|^6 \geq G_4(B, \phi)^7. \quad (\text{V.69.1})$$

The division-free form also behaves cleanly for $B = \emptyset$. For nonempty B , it is exactly the $r = 8$ case of Proposition II.59.1.

V.69.3 The all-moduli interface is a real difference

In the inspected statement, `lemma_9_3` quantifies over $N : \text{Nat}$ with `[NeZero N]`, without a primality assumption. The new cyclic theorem is for prime N . It is a quantitative strengthening in the paper’s prime setting, but not a proof of the repository proposition with only its numeric bound changed. A safe addition is a separate prime-modulus theorem with an explicit characteristic and size interface. The old theorem need not be overwritten.

V.69.4 A staged formalization plan

A useful dependency order is as follows.

1. Expose the exact graph second moment and prove (V.69.1); then generalize the moment inequality in finite abelian groups.
2. Define the finite Fourier coefficients of K_L , prove their nonnegativity, establish $0 \leq K_L \leq 1$, and prove the finite coefficient identity and the bound on $S_{m,L}$.
3. Define bounded affine degeneracy, prove the rank-two fiber count, and derive pairwise distinctness of generic tuples.
4. Formalize conditional product sampling only for distinct points, then prove the exact good/bad probabilities.
5. Prove the signed-expectation extraction lemma independently of Fourier analysis and instantiate it with the exceptional-set bound.
6. Separately formalize the augmented-column rank calculation and the fixed-characteristic system theorem.

No uncompiled proof script is presented as a certified theorem in this package. In particular, the current package contains no Lean module with an implicit or explicit admission disguised as a completed proof.

V.70 Finite audits and reproducibility

Source 17, delivered Section 10.

The package includes `verify.py` and its generated `verification_results.json`. The exact tests use integer arithmetic and rational numbers; NumPy, when available, supplies an additional independent numerical phase-average check. The random seed is 20261006.

The direct phase checks had maximum absolute discrepancy less than 2.1×10^{-17} from the exact rational answers. In the full enumeration of additive quadruples in $\mathbb{F}_2^2$, there were 64 tuples, of which 40 were affine-degenerate; the projective union bound was 48. The repeated-point example in Section II.64 is also checked exactly.

The graph tests include all nonempty partial maps from $\mathbb{F}_3$ to $\mathbb{F}_3$, deterministic random partial maps over $\mathbb{F}_5$ and $\mathbb{F}_7$, and explicit affine equality cases. Higher energies are computed by iterating exact convolution of the graph counting measure and summing squares of the resulting representation counts. This is independent of the Fourier-interpolation proof.

Check performed	Number
Exact Fejér signal lower and upper bounds	1016
Finite graphs inspected	587
Exact graph moment inequalities, orders 2 through 8	4109
Exact equality classifications, orders 3 through 8	3522
Exact generic good/bad Fourier survival identities	48
Independent direct averages over all phase parameters	16
Exhaustive scalar-linear-functional rank checks	96

Table V.70.1: Audits completed for this manuscript. Counts are script outputs.

These finite computations test normalizations, powers, equality cases, rank formulas, and collision handling. They do not certify the large parameter theorem, whose proof is the preceding sequence of lemmas. The general weighted theorem and arbitrary-system rank-gap corollary are supported by their written proofs, not by a claim of exhaustive computational verification over their infinite parameter ranges.

To reproduce the computations and build the document:

```
python3 verify.py
pdflatex -interaction=nonstopmode -halt-on-error article.tex
pdflatex -interaction=nonstopmode -halt-on-error article.tex
```

A third LaTeX pass can be used if the local installation requests an additional cross-reference update. The supplied `build.sh` performs three passes.

V.71 Further research questions and proposed directions

Source 17, delivered Section 11.

1. The actual extremal restriction exponents. For prime cyclic groups, determine the largest guaranteed size of a subset on which $1 - \eta$ of the m -term relations are respected, as a function of the absolute relation parameter α . Does one need losses of order $\alpha^m \eta^{m-1}$ in tuple count? Proposition II.62.2 settles only a scalar proof template. A matching extremal construction, or a stronger unrestricted theorem, would distinguish the method's cost from the phenomenon itself.

2. Reducing the ambient-size threshold. The union bound pays for every bounded affine degeneracy as though it survived with probability one. A rank- or multiplicity-stratified estimate of its actual survival should be tested. Can the degree $2m - 1$ in $(\alpha\eta)^{-1}$ in (F.22.5) be reduced while preserving the retention powers? This would make the present theorem substantially less asymptotic.

3. An extremal filter problem. For a prescribed mean q and degree D , maximize $\sum_n |\hat{h}(n)|^m$ among $0 \leq h \leq 1$. The Fejér filter gives a transparent feasible point, while Lemma II.62.1 gives a universal ceiling. Identifying exact or asymptotically optimal filters could improve the constants and the degree needed before controlling exceptional tuples.

4. Composite cyclic moduli. Develop a replacement for bounded affine genericity that records the Smith normal form or the relevant greatest common divisors of coefficient minors. The goal is a theorem with explicit dependence on torsion, not an unjustified carry-over of field fiber counts. Such a theorem would match the all-moduli Lean interface more directly.

[write] **7 October 2026; sources 54–59.** *Answered for one balanced relation by source 58* (Section II.276): on every cyclic group, with the domain torsion entering through $\tau_G(2D) \leq 2D$ and repeated entries through a separate collision term (its Theorem II.282.1), and with exponent 54 in place of 2^{19} for sixteen entries (Theorem A), for arbitrary abelian targets; the exponents here remain better on prime fields. Systems of relations are source 58’s Question V.219.6 (item II.8).

5. The higher arrangements of Sections 12 and 15. Compute the complete short-relation lattice for the configurations used in the later stages of Gowers’s proof, then evaluate a Fejér-type filter on that lattice. A successful extension must control every additional relation class, not only a distinguished balanced relation. This is the specific missing step before the present local improvement can be propagated through the main induction.

6. Preserve chosen spectral information while restricting. Impose auxiliary lower bounds on selected Fourier coefficients, or on Fourier energy over a prescribed set of characters, during extraction. The signed objective could be augmented by Lagrange multipliers, but one must show that its expected margin remains positive. This targets the compatibility issue that motivates the transition from Section 9 to the more delicate later arguments.

7. From approximate relations to exact structure. Under the new retention bounds, determine when a further restriction can be made an exact Freiman homomorphism of the same order without losing the prescribed spectral data. A comparison should separate the cost of obtaining a high respected proportion from the cost of making it exact. Ordinary energy extraction alone does not preserve every additional constraint needed downstream.

8. Exploit the entire error-rank distribution. The certificate (II.63.8) is a polynomial in q whose coefficients are the rank profile. Study inputs where most, but not all, unrespected tuples have high error rank. A useful theorem might combine a small exceptional low-rank tail with the improved exponents in Theorem II.63.5, rather than assuming a uniform rank gap.

9. Efficient certification and extraction. Find algorithms for estimating the signed objective and the exceptional mass without enumerating all tuples. Convolution and fast Fourier transforms can evaluate several unweighted counts, but repeated entries, conditional Bernoulli choices, and vector-valued graph data must be handled correctly. The desired output is a runtime theorem as well as an existence theorem.

10. Certified optimization in Lean. After formalizing the exact-data certificate, allow L to be chosen by a verified finite search using exact rational values of $S_{m,L}$. This would avoid freezing the convenient but nonoptimal choice $L = \lceil (m+1)/(\alpha\eta) \rceil$ in every application. The certificate is a suitable interface between computation and the structural mathematical proof.

V.72 Conclusion

Source 17, delivered Section 12.


A high-degree positive Fourier filter can replace the tiny distinguishing advantage of repeated degree-one filters. For the single balanced relation, this yields a fully explicit restriction theorem with losses $\alpha^m \eta^{m-1}$, including all repeated and degenerate tuples in the final count. Exact second-moment bookkeeping independently removes an unnecessary density factor from the input obtained from additive energy.

The fixed-characteristic analogue reveals the more general organizing quantity: the increase in rank when the graph coordinates are appended to a configuration. Its full rank-profile certificate and rank-gap corollary provide a systematic extension beyond a single equation. The remaining work is concrete: improve the treatment of exceptional tuples, understand optimal filters, and analyze the actual relation lattices of the higher configurations before making any global claim.

V.73 A compact proof-dependency map

Source 17, delivered Appendix A.

The proof of the prime-cyclic estimate depends on the following chain:



The moment inequality is independent of this chain and supplies a stronger input. The finite-field systems theorem uses the same signed objective but replaces Fourier coefficients by the rank of augmented columns. Neither proof imports a later section of Gowers’s argument.

V.74 Source provenance and claim ledger

Source 17, delivered Appendix B.

Item	Status and source
Printed target	Gowers, Lemmas 9.1–9.3 and Corollary 9.4, pp. 515–518. The original PDF pages were visually inspected; the repository’s TeX transcription was also read.
Repository snapshot	Commit <code>dca9f06383ccdfc2ed208ccb7c9882214fdb09ac</code> . The relevant statement and moment-proof source ranges were inspected. No whole-repository build was performed.

Item	Status and source
Moment inequality	Standard interpolation argument, reproved with exact mass and its unweighted equality classification. Not claimed as a new general principle.
Fejér restriction	Quantitative refinement established in this note by the displayed proof. Bibliographic priority is not established.
Weighted extension	Included in the main theorem and its proof. Its hypotheses concern weighted initial tuple counts; its conclusion is an ordinary subset.
Finite-field systems	Written proof via augmented-column rank, with explicit genericity, threshold, and error-rank hypotheses.
Method limitation	An upper bound on the scalar signed-expectation certificate only; not an impossibility theorem for general extraction.
Computations	Exact finite audits and a separate numerical phase check. No inference of universal validity solely from experiments.
Formal verification	Not performed for the new results. A dependency-level integration plan is given instead of an unverified claim of completion.
Global bounds	No new global Szemerédi bound claimed. Later arrangement analysis and a complete quantitative induction remain separate tasks.

V.75 Quantitative consequences and limits of substitution

Source 18, delivered Section 6.

The preceding results improve different numerical quantities. Keeping those quantities separate makes their mathematical significance easier to assess. Table V.75.1 records the direct consequence of each part.

Source component	Proven replacement	Direct effect
Lemma 15.4	Exact leading coefficient $k+C_d$, with explicit second-order error	A sharp asymptotic and a smaller finite exceptional-set bound
Lemma 15.5	$\lambda p \geq k + 2585445$ suffices in the stated parameter range	A smaller ambient-prime threshold; the retention exponent stays fixed
Lemma 3.10	Complete degree-four and degree-five cube coefficients	A tighter cube-count upper estimate, with all small-torsion contributions visible
Lemma 5.9	Simultaneous induction in degree through a recurrence radius	Polynomial dependence on the number of phases in the exponent, using the cited modern input
Quadratic specialization	Shared leading frequency, arbitrary lower slopes	Linear dependence on q in the reciprocal exponent, using the source’s scalar recurrence

Table V.75.1: The scope of the local improvements. Every replacement has its own hypotheses, as stated in the corresponding theorem.

V.75.1 Interpreting the polynomial exponents

The degree-two comparison is particularly transparent. For oscillation r^{-2a} , the old reciprocal exponent is 2^{11q+1} , the new general quadratic reciprocal exponent approaches $126q^2 + 129q + 2$, and the shared-leading-coefficient specialization gives $99q + 98$. The obstruction for unrestricted quadratic phases is $7q + 1$. These numbers are not four theorems with identical hypotheses: the shared frequency restricts the input class, and the obstruction is a necessary condition, not a construction. The strict inequalities in the sufficient asymptotic statements absorb coefficient-independent constants.

q	Old reciprocal	General quadratic	Shared leading term	Obstruction
1	2^{12}	257	197	8
2	2^{23}	764	296	15
4	2^{45}	2534	494	29
8	2^{89}	9098	890	57

Table V.75.2: Reciprocal exponents at oscillation r^{-2a} . The middle two columns are strict asymptotic sufficient bounds, and the last column is a necessary upper limit on a for the unrestricted class.




Figure V.75.1: Comparison of the exponents in Table V.75.2 on a logarithmic vertical scale. Larger a gives longer progressions. The figure compares exponents, not the finite input sizes at which the asymptotic statements become effective.

The main analytic reason for the improvement is visible in the recurrence $E_j = (1 + js_j)E_{j-1} + 2s_j$. For fixed degree, a recurrence cost $s_j = O_j(q)$ is paid only once at each degree. The original repeated refinement pays a degree-dependent exponent once for each polynomial. This explains the change from exponential to polynomial dependence on q . It also explains the remaining gap: a direct induction multiplies the dimension cost at successive degrees, while the necessary frontier suggests

that a more simultaneous treatment might avoid some of those products.

V.75.2 Why the cube estimate is a genuine upper-bound improvement

For $\delta, \varepsilon \geq 0$, the upper bound in Corollary III.40.9 is never larger than $(\delta + \varepsilon)^m$. The degree-one through degree-three terms have been removed; $B_d \leq \binom{m}{4}$ and $Q_d + S_d \leq \binom{m}{5}$, because their terms count disjoint classes of vertex subsets; the remaining binomial terms are the same. At fixed positive density, the first possible error is of fourth order. In the source’s uniformity parameter $\alpha = \varepsilon^{2^d}$, it is of order $\alpha^{4/2^d}$.

The sharpness theorem concerns the bound over all finite abelian groups. Its extremizing sequence lives on increasingly large elementary abelian two-groups. It does not settle the best coefficient after excluding two-torsion, nor does it classify all nearly extremizing sets. The exact formula, which separates E_2 from the ordinary fourth Fourier moment, provides information that can be used to investigate those sharper questions.

The fifth coefficient has no fixed sign. Treating it as an error requires an absolute value or a bound on its magnitude. Its explicit spectral form can be more useful when the Fourier support is known: absence of doubling relations can suppress T_5 , and absence of order-three characters removes E_3 . The remainder bound remains valid without either simplification.

V.75.3 What a full quantitative propagation would require

A new global quantitative theorem would require a closed parameter audit through the later structural and phase-removal steps. In particular, one would have to track how a changed ambient-size threshold affects every recursive call, how a new simultaneous polynomial partition is invoked inside common-difference boxes, and how the resulting widths and retained densities survive the final density iteration. An exponent improvement in one partition lemma cannot be inserted into an unchanged downstream formula without checking those dependencies.

The arrangement improvement has a particularly clean local interface: the Riesz-product proof uses the exceptional-set estimate only to absorb degenerate arrangements. Section II.66 redoes that absorption explicitly. The cube expansion is also independent of any inverse theorem. These two components can therefore be reused immediately as standalone lemmas. The polynomial transfer has an explicit analytic interface, but the general numerical exponent quoted from Lau’s preprint remains dependent on that input [73].

V.76 Relation to the Lean development

Source 18, delivered Section 7.

The repository inspection was pinned to commit

dca9f06383ccdfc2ed208ccb7c9882214fdb09ac.

The selected Lean files read during this work have matching blob identifiers in that commit. [6 October 2026; sources 15–20] The degeneracy predicate is `GeneralArrangement.IsDegenerate` (`Sections14_15.lean`, line 354, at HEAD 0b4890a1c); `lemma_15_4` and `lemma_15_5` (lines 402 and 414) have exact companions (notes after Theorem II.66.8 and Corollary II.66.9); `lemma_5_9`

(`Section05.lean`, line 261) and `lemma_3_10` (`Sections01_03.lean`, line 340) also have exact companions. None of this source’s statements is formalized. The paper source was read as blob 41085a0efde8c1932e86e80c791984f52841c22e. The accompanying provenance file records the exact paths and identifiers. No repository changes were made as part of preparing this report.

The arrangement refinement uses the existing definitions. `GeneralArrangement` stores the shared side vector, cube bases, and cross-sections. Its degeneracy predicate quantifies over all coordinate monomials, which is precisely the stronger definition used here. Thus the new proof does not need a change to the meaning of degeneracy. It needs a stronger counting theorem for that predicate.

Inspected component	Role in a possible implementation
<code>Sections14_15.lean</code>	Supplies <code>GeneralArrangement</code> , <code>arrangementMoment</code> , the full degeneracy predicate, and the current statements of Lemmas 15.4 and 15.5.
<code>Proofs15DegenerateArrangements.lean</code>	The current proof separates zero sides from coefficient cases. The refined proof would replace its single-polynomial fibre estimate with the paired affine-rank bound.
<code>Proofs15Walsh.lean</code>	A starting point for the Boolean moment transform and the characterization of parity by vanishing proper moments.
<code>Proofs15LevelSets.lean</code>	Proves the general multilinear level-set estimate. The refined arrangement argument bypasses this estimate for nonparity certificates by exhibiting explicit affine minors.
<code>Sections01_03.lean</code>	Records the uniformity and cube-count statements whose normalizations must be respected.
<code>Proofs03CubeUpper.lean</code>	The current Lemma 3.10 proof uses Minkowski’s inequality. The new expansion would be a separate stronger theorem proved from finite sums, Fourier orthogonality, and cube Cauchy–Schwarz.
<code>Section05.lean</code>	Contains the polynomial-partition statements. A recurrence-radius premise provides a clean interface for checking the elementary degree-descent argument independently of its analytic input.

For cube norms, the article uses normalized averages. The source’s unnormalized quantity differs by the factor $|G|^{(d+1)/2^d}$. Equivalently, the 2^d -th power of the source norm is $|G|^{d+1}$ times the normalized cube moment. Recording that equality once would prevent repeated normalization errors in later statements.

An economical implementation order is: first the Boolean moment inversion and the two affine-rank counting lemmas; then the single-certificate bound; then certificate equivalence and its finite class count; finally the sharp hyperplane count and the absorption corollary. A simpler finite theorem with a larger coefficient could be checked before the exact combinatorial count C_d is implemented. In the cube part, the universal finite Fourier kernel identity should precede the classification of small matrices. For partitions, the two elementary integer decomposition lemmas and the degree-descent transfer can be formalized with recurrence as a named hypothesis.

The new theorems in this report are mathematical proofs, not newly checked Lean declarations. The inspected source establishes the integration points; it does not certify the results written here. No Lean compilation of these additions is claimed.

V.77 Exact finite verification and reproducibility

Source 18, delivered Section 8.

The accompanying `verify_refinements.py` uses only the Python standard library. It checks finite consequences of the proofs with exact integer and rational arithmetic. The program is corroborating evidence and a reproducible way to inspect the constants; the general arguments above remain the proofs.

The following checks passed:

1. All $\binom{16}{4} + \binom{16}{5} = 6188$ four- and five-vertex subsets of the four-dimensional cube were classified by rank over suitable finite fields. The counts were 100 ordinary rectangles, 40 extra binary planes, 160 ordinary five-circuits, and 16 extra ternary configurations.
2. Fourth and fifth coefficients were computed directly by expanding the product over every cube, for dimensions three and four on cyclic groups of orders 2, 3, 4, 5, 6, 7. All twelve test cases agreed with the formula, including groups with two-torsion, three-torsion, both, and neither.
3. Pure-parity hyperplanes were enumerated independently for $d = 2$ at $p = 5, 7$, for $d = 3$ at $p = 7, 11$, and for $d = 4$ at $p = 11, 101$. The resulting counts were $C_2 = 4$, $C_3 = 51$, and $C_4 = 484$, as predicted.
4. Ternary certificate classes modulo affine equivalence were enumerated at sizes 2, 4, 6, 8, giving 1, 25, 301, 3025. These agree with $(3^n - 2^{n+1} + 1)/2$.
5. The exceptional-diagonal probability for the sharp nonparity certificate was computed exactly for $d = 2$ and $p = 5, 7, 11$. The values $29/625$, $55/2401$, and $131/14641$ agree with $p^{-2} + (p - 1)p^{-4}$.
6. Every binary square matrix of sizes two, three, and four was enumerated. The maximum absolute determinants were 1, 2, 3, consistent with the independent Hadamard-bound proof.

As a concrete check of the order-three correction, on $\mathbb{Z}/3\mathbb{Z}$ take $f = (2, -1, -1)$ and cube dimension four. The fourth coefficient is 200 and the fifth is 352. Omitting the sixteen ternary configurations would give a fifth coefficient of 320, so this test detects a real failure of the torsion-free formula in characteristic three.

To reproduce the checks, run

```
python3 verify_refinements.py --output verification_results.json
```

To rebuild the article, run `pdflatex` twice on the supplied main `.tex` file, keeping `partition_exponents.pdf` beside it. The supplied source is self-contained apart from that included figure and standard L^AT_EX packages; its bibliography is embedded.

V.78 Further research questions and proposed next steps

Source 18, delivered Section 9.

The following questions arise directly from the proofs. They distinguish well-defined improvements to these local statements from the much larger task of improving global Szemerédi bounds.

1. **The second degeneracy coefficient.** Determine the coefficient of p^{-2} in the count of degenerate d -arrangements, initially for $d = 2$ and small k . The first-order proof shows what must be included: intersections of the leading hyperplanes and the nonparity loci of codimension two. A satisfactory answer should identify any exceptional primes, rather than fitting a polynomial to a few counts.

[write] 6 October 2026; sources 21–27. *Answered for $d = 2$* by source 27’s Theorem F.32.3: $(k^2 + 21k - 6)/2$ for characteristic $p > 4 \cdot 2^k$, with the exact count for $k = 1$ in characteristic at least 11 (exceptional characteristics 2, 3, 5, 7); general d is reduced to the statistic $b_{2,d,p}$ (its Corollary II.99.2). Item II.38.

2. **Distinct nonparity loci.** The coefficient $A_{2d2^k} - A_{2d}$ counts certificate classes, including many with impossible or identical solution sets. Classify the simplest additional equivalences between the resulting algebraic loci. Even a systematic bound for the certificates of lowest moment order $k - 1$ would improve the explicit second-order error.
3. **The full moment-order distribution.** Use the parameter $t = k - m$ in Lemma II.66.2 instead of replacing all costs by their $t = 1$ upper bound. This requires counting ternary functions with specified vanishing Boolean moments. A first milestone is an exact enumeration for $k \leq 4$, followed by a structural inequality valid for general k .
4. **The balanced-subset hyperplane arrangement.** Determine its intersection lattice and characteristic polynomial. The number of hyperplanes is already C_d ; understanding their intersections would give the exact pure-parity degeneracy probability for all sufficiently large primes and would contribute to the preceding second-order problem.
5. **Small primes and one-arrangements.** Treat $p \leq 2d$ by counting balanced ternary vectors modulo p , and treat $d = 1$ separately. The diagonal cross-section event then has different size, so the p^{-2} conclusion for nonparity certificates cannot simply be retained. In characteristic three, affine symmetries of the ternary alphabet also change the quotient count.

[write] 6 October 2026; sources 21–27. *Partly answered.* For $d = 1$ the leading coefficient is $(k^2 + 3k)/p$ (source 22’s Theorem II.85.1); the first-order count for small p is source 27’s Theorem II.96.3, with $L_{d,p}$ from its Theorem II.95.1, including characteristic three. Item II.38.

6. **Sixth and higher cube coefficients.** Extend the small-matrix classification to six vertices. New affine dependence dimensions and new integer invariant factors must be tracked; a list of rational circuits alone will not suffice over all finite groups. The desired output is an exact formula with a remainder bound, not only a numerical enumeration of one cube dimension.
7. **Sharp coefficients on odd-order groups.** Determine the best fourth-order coefficient when order-two characters are absent, especially in dimension four. The universal B_d sharpness construction does not answer this question. A likely useful subproblem is an optimal comparison between a real mean-zero function’s U^2 and U^d norms under the relevant group restrictions.
8. **Stability of cube-count sharpness.** If a sequence nearly attains B_d , what must its low-order Fourier spectrum look like? The separate nonnegative terms V_4 and E_2 make this a more structured problem than the corresponding triangle-inequality equality case. Distinguish weighted functions from indicator sets in any stability statement.
9. **The general quadratic partition gap.** Can the sufficient reciprocal exponent be reduced from order q^2 to order q ? The obstruction is $7q + 1$ at error exponent $2a$. One should first decide whether flexible lengths and the stronger $v - 1, v$ requirement have the same asymptotic frontier.

10. **Eliminate several degrees in one step.** The transfer handles the leading coefficients simultaneously, but then repeats the process for lower degrees. Find a recurrence formulation that also controls the coefficients created by translation. The shared-leading quadratic corollary supplies a useful test family in which the number of genuinely independent top frequencies is much smaller than the number of polynomials.
11. **Coefficient-rank profiles and explicit constants.** Develop a theorem in which each degree has its own bounded integer span of frequencies, and compute the resulting sequence of recurrence costs. In parallel, extract fully numerical recurrence constants where possible. An exponent with an uncontrolled starting threshold is valuable asymptotically but insufficient for concrete finite comparisons.
12. **A complete formal and quantitative substitution.** Implement the elementary lemmas in Lean, then trace their use through the remaining proof with explicit parameter interfaces. The first deliverable should be a checked replacement for the Lemma 15.5 threshold and a checked conditional polynomial-partition transfer. A global quantitative claim should follow only after every later width, density, and ambient-size condition has been verified.

V.79 Source record and attribution

Source 18, delivered Appendix A.

The research request identified two paths in <https://github.com/VladimirReshetnikov/ProveIt>:

```
Combinatorics/Ramsey/Papers/sz-thm-gowers-proof
Combinatorics/Ramsey/Lean/GowersSzemerédi
```

The published PDF was also consulted through the University of Maryland mirror listed in the bibliography. The paper reconstruction is useful for searching and for seeing explicit source corrections; the article does not silently identify all edited formulations with the printed 2001 text.

The core derivations were cross-reviewed independently within this task. The analytic inputs are credited where used, and the computational files record only finite checks. This division is deliberate: a numerical experiment cannot prove the general rank theorem, and an uncompiled Lean sketch would not constitute kernel verification.

V.80 Source audit and the boundary of the Lean integration

Source 19, delivered Section 7.

The mathematical improvements in this article concern a finite sampling and interpolation step in the proof of Gowers’s Lemma 16.10. The following audit distinguishes the original published paper, the edited transcription in the repository, the propositions currently stated in Lean, and their available proofs. The repository snapshot is commit `dca9f06383ccd2ed208ccb7c9882214fdb09ac`. The source checks below are read-only inspections at that commit; a new Lean build was not run for this article.

V.80.1 Published argument versus the edited transcription

In the original paper, Lemma 16.10 and its proof occupy printed pages 574–575. The relevant finite step starts with, in each fibre, at most q classes on which the function is affine in the last coordinate. The proof uses $r = \lceil q\sigma^{-2} \rceil$ random coordinates and seeks two distinct sampled points in each retained class [1]. The edited repository transcription places Lemma 16.10 at line 3459 and retains this sampling step, while adding warnings about the closing parameter comparison [89].

There is a precise distinction between two sample hits and two distinct sampled coordinates. For independent sampling with replacement from an n -point set, a class of size s has fewer than two distinct sampled coordinates with probability

$$s \left(1 - \frac{s-1}{n}\right)^r - (s-1) \left(1 - \frac{s}{n}\right)^r, \quad s \geq 1. \quad (\text{V.80.1})$$

Indeed, the event consists of no hit, together with the disjoint events that the unique sampled member of the class is each of its s possible points. In contrast, the binomial expression $(1-u)^r + ru(1-u)^{r-1}$ bounds fewer than two hits counted with multiplicity. It does not, without another argument, bound the event in (V.80.1). A singleton class always fails the requirement of two distinct points, regardless of the sample size.

For a uniformly chosen r -element subset, where $0 \leq r \leq n$, the corresponding probability is instead

$$\frac{\binom{n-s}{r} + s \binom{n-s}{r-1}}{\binom{n}{r}}. \quad (\text{V.80.2})$$

The conventions are that an out-of-range binomial coefficient is zero and that the population has $n \geq 1$ points. The sampling theorem of the present article uses this finite-population model and the more general recovery condition $|R \cap C| \geq \min(d, |C|)$. The latter allows a fully sampled small class to be represented by a lower-degree interpolant. Thus neither repeated coordinates nor singleton classes are silently treated as a pair of distinct interpolation nodes.

V.80.2 The closing graph budget: two different source formulas

It is essential to retain a dimensional distinction between the two versions of the argument. Put

$$A_j = 2^{2^{j+8}}, \quad \mathcal{Q}(t, \gamma, j) = (\gamma t)^{-A_j}.$$

These are the graph-count controls used in the repository's `multipleC` and `multipleQ`, and agree with the displayed controls on printed page 567 of the original paper [1, 90]. Write

$$\sigma = \rho/4, \quad R = \gamma^{-2} s(2^{-(k+2)} \theta, \gamma, k), \quad p = 4r^2 R,$$

where r is the number of sampled coordinates. The original published comparison on page 575 is

$$\mathcal{Q}(\sigma/p, \gamma, k)^p \leq \mathcal{Q}(\rho, \gamma, k+1). \quad (\text{V.80.3})$$

The edited repository transcription instead uses dimension $k+1$ on the left-hand side as well [89]. These are different inequalities and should not be quoted interchangeably. [6 October 2026; sources 15–20] The printed page was not checked at the write; this is the source's reading of the 2001 paper. The transcription carries an editorial note at Lemma 16.10 that the displayed closing comparison does not justify the lemma.

Both versions have an explicit algebraic obstruction with the displayed controls. Assume $k \geq 1$, $r \geq 1$, and $0 < \theta, \gamma, \rho \leq 1$, and set $a = \log(1/(\gamma\rho)) \geq 0$. Since $A_{k+1} = A_k^2$, taking logarithms shows that (V.80.3) is equivalent to

$$p(a + \log(4p)) \leq A_k a. \quad (\text{V.80.4})$$

In particular it fails whenever $p \geq A_k$. The displayed auxiliary parameter already exceeds that range. Indeed, with $B_k = 2^{2^{k+6}}$, the definition of s gives

$$R = \gamma^{-2} \left(\frac{2^{k+3}}{\theta\gamma} \right)^{B_k} \geq 2^{(k+3)B_k} > 2^{2^{k+8}} = A_k.$$

For the last strict inequality it suffices to observe that $B_k > 2^{k+8}$ for $k \geq 1$. Therefore $p = 4r^2 R > A_k$, and (V.80.4) fails throughout this range. This calculation concerns the stated numerical comparison, not the truth or falsity of Lemma 16.10 as a mathematical assertion.

The same-dimension comparison in the edited transcription fails even more directly: for $p \geq 1$,

$$\mathcal{Q}(\sigma/p, \gamma, k+1)^p = \left(\frac{4p}{\gamma\rho} \right)^{A_{k+1}p} > \left(\frac{1}{\gamma\rho} \right)^{A_{k+1}} = \mathcal{Q}(\rho, \gamma, k+1).$$

[write] **6 October 2026; sources 15–20.** *Certified formally since the pin:* `section16_same_dimension_closing_comparison_fails` (Proofs16ClosingComparison.lean, line 25, HEAD 0b4890a1c) proves $\text{multipleQ}(\rho, \gamma, k) < \text{multipleQ}(\rho/(4p), \gamma, k)^p$ for $p \geq 1$ and $0 < \gamma, \rho \leq 1$, this display in the Lean indexing; like this source, it is an obstruction to the displayed proof step, not a counterexample to Lemma 16.10.

Consequently, reducing the number of anchors gives a real local improvement but does not by itself justify either displayed closing comparison. One must also redesign the quantitative closure, preserve the dependence on the outer parameter, or replace the relevant structural input. The present article does not infer an improved global Szemerédi bound from the sampling improvement.

V.80.3 What the pinned Lean development establishes

At the pinned commit the machine-readable ledger contains all 120 numbered catalogue entries and records 98 exact companion proofs and 22 open statements. The repository checkpoint reports that its 169 reachable modules compiled and that its axiom audit checked 1,658 public theorems, allowing only `propext`, `Classical.choice`, and `Quot.sound` [91, 92]. [6 October 2026; sources 15–20] *Stale since the pin:* at HEAD 0b4890a1c the ledger records 101 exact companions and 19 open statements and `FORMALIZATION_STATUS.txt` 233 modules and 2,138 audited public theorems; `lemma_16_9`, `lemma_13_8` and `lemma_13_7` have closed since dca9f0638. These are the results reported by that checkpoint; they are not a new build result of this article. The audit itself explicitly distinguishes an axiom check of a conditional implication from a proof of its hypotheses.

The relevant part of the catalogue is as follows.

Statement	Status at the pinned commit
Lemma 16.3, the one-dimensional base case	Exact companion
Lemmas 16.5 and 16.6	Exact companions
Lemma 16.8, with repaired parameters	Exact companion
Theorem 16.2 and Lemma 16.4	Open
Lemmas 16.9 and 16.10	Open
Corollary 16.11	Open
Theorems 18.1 and 18.2	Open

[write] 6 October 2026; sources 15–20. *Stale row.* Lemma 16.9 is now an exact companion: `lemma_16_9` (`Section16.lean`, line 749) is proved by `lemma_16_9_holds` (`Proofs16Lemma9.lean`, line 158), added in 16bbbe5ad at 14:31 on 6 October 2026, after this source’s pin (14:17). Lemma 16.10 is still open (`lemma_16_10`, line 851), still with no related theorem in the ledger, although proved partial results exist (note after Corollary II.70.8).

In particular, `lemma_16_10` has no exact companion or related theorem in the ledger. Its two substantive inputs are stated as `FinalCoordinateSectionsMultiplyLinear` and `Section16AllBoxLineCovers` [90]. The first gives multiple multilinearity of each section at a fixed final coordinate, with iteration parameter

$$\gamma^{-2}s(2^{-(k+2)}\theta, \gamma, k).$$

The second packages the line covers and width bounds intended to come from Lemma 16.9.

Several interface facts matter for an implementation. `Section16LineCover` stores at most a natural number `qGamma` of affine functions in each fibre. The separate natural number `qDelta` controls the recurrence-width loss. They have different bounds and cannot be identified without a proof. The migrated `MultiplyLinear` predicate requires proper input and output boxes and a loss parameter in $(0, 1]$. In contrast, the current `Section16AllBoxLineCovers` interface still quantifies over arbitrary boxes and does not yet carry the corresponding properness clauses. [6 October 2026; sources 15–20] *Stale:* at HEAD 0b4890a1c `Section16AllBoxLineCovers` (`Section16.lean`, line 824) requires `P.IsProper`; the older unrestricted form is kept as `Section16AllBoxLineCovers_unrestricted_box_encoding` (line 806). The current signature of `lemma_16_10` also assumes only a nonzero modulus. Distinct interpolation nodes in an arbitrary `ZMod N` need not have invertible differences. A field hypothesis, a prime modulus, or explicit pairwise unit-difference hypotheses are therefore needed in the algebraic interpolation layer.

V.80.4 A practical formalization route

The finite sampling theorem should first be formalized independently of the Section 16 induction. Its natural data are a finite population, a finite family of rows, at most q disjoint classes in each row, and the uniform family of r -element population subsets. Counting pairs consisting of a sample and a class gives the hypergeometric expectation. Finite averaging then supplies an actual sample with the claimed lost mass. This stage needs finite sums and binomial coefficients; it does not need a probability-space formalization or a prime modulus.

The next layer constructs the Lagrange interpolant for each nonempty anchor subset of size at most d over a field. For a recovered large class, agreement at d nodes implies agreement with its degree at most $d - 1$ polynomial. For a smaller recovered class, all its points are nodes, which is sufficient for agreement on that class. The interpolation statement should expose both conclusions rather than assert uniqueness of the original polynomial on a singleton.

Finally, the sampled fixed-coordinate sections must be placed on a common proper refinement, and their covering functions substituted into the interpolation formula. The existing public helpers `BaseCase.properMultiplyLinear_finsetUnion` and `BaseCase.properMultiplyLinear_fin_sum` provide proved finite-family closure for iteration parameters at least one; `IsMultilinear.mul_new_coordinate` supplies the elementary multilinear lift used in the affine case. General degree in the final variable requires an additional mixed-degree predicate and its closure lemmas. Domain restriction also matters: a sampled section is only certified where that section’s partial domain is present.

The resulting local theorem can be connected to `Section16LineCover`. It should retain explicit sample, graph-count, width, and exceptional-mass bounds. Declaring `lemma_16_10_holds` would require the additional geometric and quantitative closure obligations described above. No such full Lean companion, and no completion of Theorems 18.1–18.2, is claimed here.

V.81 Verification and reproducibility

Source 19, delivered Section 9.

The package contains a standalone L^AT_EX source with an embedded bibliography, the compiled PDF, the figure used in the article, the source audit, and executable diagnostics. The mathematical proofs do not depend on a numerical search finding the claimed objects. Computation is used to check exact finite identities, normalization, endpoint conventions, and illustrative values.

V.81.1 Exact finite checks

The script `verification/verify_sampling.py` uses only Python’s standard library and exact rational arithmetic. Its recorded run checks the following.

- 1,136 comparisons between the hypergeometric formula and explicit enumeration of subsets, for populations of size at most eight and interpolation orders at most four.
- 1,664 instances of the universal loss bound, including the finite dynamic program and the full-population endpoint.
- All 64 subsets of a six-point population in the symmetric extremal-family construction, grouped by their sample sizes. The loss is independent of the particular anchor subset and equals the exact minimax value.
- A deterministic conditional-expectation selection example with 25 partial rows on 12 columns. Its five selected columns are 2, 5, 9, 10, 11, with zero-based column indices. The expected unnormalized loss decreases from 21047/198 to an actual loss of 78 out of 300 ambient row-column points.
- 243 finite-field interpolation evaluations, including small classes and functions with independent affine degree bounds in two horizontal variables.

The file `verification/sampling_results.json` records the results and the exact profiles in table II.68.1.

The script `verification/verify_polynomial.py` checks the determinants, the finite Newton-lift identities, and the modular counterexample described in section I.34. Its exhaustive quadratic tests examine 20,402 coefficient tuples, and its determinant checks cover 40 degree-spacing pairs. Results are in `verification/polynomial_results.json`.

The figure script `verification/make_figures.py` evaluates the unique-root equations for the occupancy constants with high-precision decimal arithmetic and draws the figure using Matplotlib. The displayed decimals are numerical evaluations, not a replacement for the exact formulas and uniqueness proof.

V.81.2 Build and source provenance

From the unpacked directory, the PDF can be rebuilt with

```
latexmk -pdf -interaction=nonstopmode -halt-on-error \  
gowers_interpolation.tex
```

The supplied PDF figure is sufficient for compilation; regenerating it is optional. The file `README.txt` gives the short commands for rerunning the diagnostics.

The source manifest records the exact repository commit and file hashes used for comparison. The full paper and Lean development were read through the selected GitHub integration. The relevant published formulas were also checked against the original PDF, including its displayed Section 16 controls and final graph comparison. The source audit deliberately records the original and edited formulas separately.

All principal new claims received a written-proof review independent of their initial derivation. This is a description of the checks performed for the report. It is not external peer review, and no new Lean build or kernel verification was run.

V.82 Further research questions and proposed projects

Source 19, delivered Section 10.

The questions below are proposed follow-up problems arising from the proved results. Unless an external source is cited, they are questions formulated here, not assertions that a particular problem is already recognized as open throughout the literature.

V.82.1 Exact sampling and interpolation

Research question V.82.1 (Structure of finite extremal profiles). Can the maximizers in (II.68.1) be classified without dynamic programming? In the numerical examples, equal small classes often maximize the loss, but the small-sample regime also uses unequal sizes forced by the total-size constraint. Determine how many distinct positive class sizes a maximizer ever needs, and locate the transitions as r/n varies.

The exact integer optimization makes this a finite mathematical question with certificates. A useful first result would be a uniform bound, depending only on d , on the number of distinct sizes in an optimizing vector. Such a result would turn the extremal value into a substantially smaller optimization.

Research question V.82.2 (Rates in the sharp sampling law). Give explicit upper and lower error bounds for $r\mathcal{M}_{n,q,d}(r)/q - C_d$ in terms of $1/r$, r/n , and q/r , with a stated range of validity. Can these bounds replace the simple factor d in the finite sample budget by a nearly optimal constant at a practically verifiable population size?

The proof of theorem II.69.4 already separates compact Poisson approximation from an exponentially controlled tail. The remaining work is quantitative uniformity and the lattice rounding of the maximizing class size.

Research question V.82.3 (Selecting cheap and informative slices). Suppose different columns have different structural refinement costs. Can the anchor selection minimize a combination of lost mass and the elementary-symmetric graph count, rather than the lost mass alone? What is the correct replacement for uniform sampling in that problem?

Here a column should be assessed by both the classes it helps recover and the complexity of its slice cover. The heterogeneous formula in remark II.70.4 supplies a precise objective; the conditional-expectation algorithm supplies a starting point when that objective has a computable expectation.

Research question V.82.4 (Going beyond the common-anchor certificate). Can algebraic relations among the row polynomials reduce the number of required sampled slices below the q/σ scale? For example, what changes when many affine pieces share a slope, or when all slice functions belong to a low-dimensional common function space?

The minimax lower bound does not forbid such an improvement: its rows deliberately have no structure beyond the specified class partitions. A theorem using additional shared algebraic information would address a different, stronger input model.

Research question V.82.5 (Robust polynomial lifting). Develop an analogue when each row is approximately polynomial or each sampled slice has errors in its values. Which node separation, agreement-density, or norm hypotheses permit a quantitative repair without losing the linear number of slice refinements?

Over the reals or complex numbers, the norms of Lagrange basis polynomials must be controlled. Over finite fields, an agreement-based formulation may be more appropriate. These are different problems and should have separate hypotheses.

V.82.2 Geometry and polynomial recurrence

Research question V.82.6 (Proper modular boxes without inherited coordinates). Replace the additional inherited-coordinate hypothesis in corollary II.70.8 by the precise proper-box predicates used in the live Lean development. Prove a synchronization theorem with an explicit width loss, or produce an obstruction showing what additional information is necessary.

This is the most immediate geometric integration task. The algebraic interpolation theorem is already independent of it. A proof should explicitly control subprogression steps and wrapping, rather than infer integer containment from injectivity of a modular parameterization.

Research question V.82.7 (The nonlinear exponent gap). For fixed $k \geq 2$, can the sufficient partition bound be improved from a denominator of order q^k to a smaller power of q ? Is the necessary region

$$\left(1 + \sum_j j q_j\right) a + \left(\sum_j q_j\right) b \leq 1$$

attainable, even for quadratic phases?

An intermediate improvement by one power of q would already require a different way of handling lower coefficients created by translation. The present proof repeatedly solves a new lower-degree recurrence problem; retaining correlations between those problems is a plausible place to seek a gain.

Research question V.82.8 (Prescribed lengths without the square loss). Can a complete partition into cells of two consecutive prescribed lengths achieve the flexible-length exponent range without the generic square-length requirement in lemma I.34.4? What is the optimal tradeoff among minimum length, allowed length variation, and phase diameter for nonlinear phases?

Research question V.82.9 (Numerical recurrence thresholds). Make the coefficient-uniform constants $U_{j,m}$ in the recurrence input numerical, at least for small degrees and moderate dimension. Track them through (I.34.9) before attempting any fully numerical global density-increment estimate.

The numerical exponent specialization from Lau does not by itself supply numerical values for every threshold. A complete parameter audit must account for both.

V.82.3 Closing the structural induction

Research question V.82.10 (A consistent family of structural controls). Design control functions for the induction that retain the outer parameter dependence and absorb the actual output of the shared-slice theorem. Can one prove a closure statement using those controls and then quantify the resulting final density increment?

The logarithmic obstruction in section V.80.2 provides an immediate test for any proposed repair. A revised definition must be stable under the operations actually used: restriction, union, sums, adding an unused coordinate, choosing slices, and interpolation. Passing one numerical comparison is not enough if a later operation violates the same control family.

Research question V.82.11 (Formal certificates for the finite ingredients). Formalize, in order, the hypergeometric counting identity, the marked-sample bound, the finite minimax construction, and the common-refinement interpolation theorem. Can the exact dynamic-programming profiles be exported as small certificates checked independently of the search program?

These tasks have clear finite statements and relatively modest dependencies. They can be completed before the full Section 16 induction is repaired. The higher-degree phase-partition transfer can likewise be formalized against an explicit recurrence hypothesis, keeping the imported theorem separate from the new transfer argument.

Research question V.82.12 (A complete local-to-global parameter audit). After the geometric and structural closure tasks are complete, propagate all changed constants through the remaining inverse, phase-removal, density-selection, and iteration steps. Which local savings survive in the final function of density and progression length?

This last task is the necessary route from the improvements in this article to any claim about a global Szemerédi bound. The present report supplies exact local theorems, sharpness certificates for a natural sampling problem, and explicit interfaces for carrying out that audit.

V.83 Formalization and reproducibility

Source 20, delivered Section 9.

V.83.1 Connection with the inspected Lean interfaces

The inspected `Sections01_03.lean` defines `progressionAverage`, `theorem_3_2`, `corollary_3_3`, `lemma_3_8`, and `lemma_3_9` [94]. The progression average there is an unnormalized sum on $\mathbb{Z}/N\mathbb{Z}$. Its Theorem 3.2 specification assumes prime modulus and $k \leq N$ [6 October 2026; sources 15–20; *HEAD 0b4890a1c*] (`theorem_3_2`, `Sections01_03.lean`, line 258, and `corollary_3_3`, line 276, with exact companions `theorem_3_2_holds` and `corollary_3_3_holds`, `Proofs03ProgressionAverage.lean`, line 381, `Proofs03ProgressionCount.lean`, line 349; none of this source’s results is formal); the new cyclic statement can instead use (III.44.4), or impose the exact coprimality condition to make all costs one. The inspected `Proofs01_03.lean` contains the finite Fourier identities underlying this normalization [95]. These are source observations, not a claim that the repository already formalizes the theorems in this article.

A useful implementation order is as follows. First formalize uniform fiber averaging for a surjective homomorphism, including the equal-fiber cardinality proof. Next formalize the joint-image subgroup and its two quotient maps from Lemma III.43.1. Then define directional cube powers as normalized finite sums and establish positivity without real fractional powers. The kernel-elimination theorem should first be proved in its powered form (III.43.6); the final roots are a separate real-analysis wrapper. The finite-index comparison can similarly be proved as an inequality between 2^s -th powers using character orthogonality and the existing cube Cauchy–Schwarz interface.

The second-order telescoping identity is independent of groups and analysis. It is a finite polynomial identity over a commutative ring, with the first two selected indices providing a natural induction or finite-set partition proof. Separating it from the analytic bound should make both developments reusable. The Gauss construction requires a separate finite-field exponential-sum development and should not block formalization of the main counting theorem.

No Lean executable was available in the working environment used for this package. No Lean source is presented as compiled or kernel checked, and the repository was not modified. The supplied verification program uses exact rational arithmetic for finite tests, not proof-assistant certification of the general statements.

V.83.2 Executed verification

The accompanying `verify.py` uses the Python standard library for its exact tests. It computes directional cube powers by derivative recursion. For the two-function theorem it compares rational quantities in the powered inequality

$$|\Lambda|^{2^{k-1}} \leq D_i(f_i)^{2^{k-1}} D_j(f_j)^{2^{k-1}} (\mathbb{E}|f_\ell|^2)^{2^{k-2}} \prod_{r \notin \{i,j,\ell\}} \|f_r\|_\infty^{2^{k-1}}. \quad (\text{V.83.1})$$

No floating-point root is needed. This tests the actual L^2 gain, including non-coprime and repeated slopes.

An optional NumPy check, also executed, directly averages the weighted three-term example on $\mathbb{F}_{29}^2$. It gives

$$\Lambda_3(f) = 0.00003419510332526960,$$

with discrepancy less than 1.4×10^{-20} from the evaluated exact formula. Its computed U^2 fourth power is approximately $7.85288905 \times 10^{-7}$, below the proved bound $4.64476813 \times 10^{-6}$. These are floating-point checks and are labeled separately in `verification_report.json`. Neither the Gaussian computation nor the finite rational tests replace the proofs.

Exact test family	Number
Exhaustive two-function/ L^2 inequalities on $\mathbb{Z}/2\mathbb{Z}$	6,561
Random rational two-function/ L^2 inequalities, orders 3–5	1,710
Finite-index norm comparisons	252
Exact equalities for coset-supported functions	252
First-two-residual polynomial identities	700
Exhaustive three-term indicator bounds on odd cyclic groups	680
Higher-order rational quadratic counting bounds	105
Quadratic-phase matrix rank cases	64
Exact pair-quotient averaging identities	65
Total	10,389

Table V.83.1: Executed exact arithmetic tests. All passed. Random tests use the fixed seed 20261006; “exhaustive” refers only to the specified finite families, not to all functions or groups.

To reproduce the exact checks, run

```
python3 -S verify.py --output verification_report.json
```

To include the optional NumPy calculation, omit `-S` and add `-gauss`. The manuscript is self-contained as a TeX source with an inline bibliography and no external figures. Three runs of `pdflatex` resolve the contents and cross-references.

V.84 Further research questions and topics

Source 20, delivered Section 10.

The questions below are a proposed program arising from the proved bounds. Except for the explicit local extremal questions formulated here, they are research directions rather than assertions that a literature-wide open status has been established.

Question V.84.1 (Optimal exponent for longer symmetric counts). *For fixed $k \geq 4$ and density bounded away from zero and one, determine the largest θ for which an estimate $|\Lambda_k(1_A) - \delta^k| \leq C_{k,\delta} U^{k-1} (1_A - \delta)^\theta$ holds uniformly over admissible finite abelian groups. The present article proves $\theta \geq 2$ and settles optimality only for $k = 3$. Can higher-degree phase rank degeneracies furnish matching examples?*

Question V.84.2 (Sharp three-term density coefficient). *Determine the best coefficient multiplying u^2 in the small- u three-term problem, separately for weights and indicator sets. Is the upper coefficient $\sqrt{\delta(1-\delta)}$ attainable asymptotically under the indicator constraint? The rounding proof preserves a positive coefficient, but is far from an optimal-constant calculation.*

Question V.84.3 (Prime-cyclic sharpness). *Construct a quantitatively controlled prime-cyclic analogue of Theorem III.48.3, or identify an obstruction that improves the available exponent or coefficient in that restricted family. A transfer argument must preserve the count excess relative to the square of the norm, rather than merely preserve positivity.*

Question V.84.4 (The best multilinear torsion constant). *For fixed L and marked positions i, j, ℓ , determine the operator norm in (III.43.5) after replacing both directional seminorms by global ones. The product $\kappa_i \kappa_j$ is explicit, but sharp individual norm conversion does not settle simultaneous sharpness of the full counting estimate.*

Question V.84.5 (Equality and stability of the two-function argument). *Classify equality in Theorem III.43.3 beyond constant functions, and obtain a stability statement for near equality. The proof suggests compatibility conditions between almost constant fiber magnitudes and almost proportional quotient-conditional derivatives. Formulating the correct quantitative metric is part of the problem.*

Question V.84.6 (Optimized unequal-density elimination). *Develop an exact finite optimization procedure for second-order elimination using the full profiles $\|b_i\|_{U^s}$, densities, L^2 norms, and subgroup costs. The existing first-order dynamic program in the repository is not automatically optimal for this two-marked problem. Can one obtain a useful polynomial-time approximation without enumerating all position orders and tail decompositions?*

Question V.84.7 (More than one ordinary norm gain). *Which systems admit an improvement retaining two uniformity norms and several L^p factors? Ordinary iteration tends to change the power of a later density gain. A sharp mixed-norm region, stated as a finite polytope of exponents when possible, would clarify the best sparse-set bounds obtainable from this method.*

Question V.84.8 (A quotient-adapted counting baseline). *Replace the product of means by a baseline determined by conditional expectations on the torsion quotients, with (III.46.3) as its first correction. Can a systematic expansion remove large κ costs while retaining global uniformity control on the remaining genuinely transverse directions?*

Question V.84.9 (Localization preserving two uniform factors). *Find a no-wrap or interval-localization mechanism that retains two small balanced norms in every error term. This would connect the quadratic modular threshold to integer density-increment arguments without assuming that multiplication by an arbitrary interval indicator preserves uniformity.*

Question V.84.10 (Effect on a complete quantitative iteration). *Insert the improved counting threshold into a fully specified inverse theorem and density-increment scheme, and propagate every parameter. Which exponent in the final bound changes, and which loss remains dominant? This is the required calculation before any global-bound improvement can be claimed.*

Question V.84.11 (General coefficient modules). *For systems of integer matrix forms on finite modules, express the subgroups $L_i(\ker L_j)$ through Smith normal forms or determinantal ideals. An exact arithmetic interface could turn the abstract theorem into a computable torsion-sensitive bound for higher-dimensional configurations, including systems with several parameter directions.*

Question V.84.12 (Kernel-checked proof certificates). *Formalize the two-function theorem, the sharp finite-index comparison, and the second-order telescoping identity as independent reusable components. An especially useful first milestone is a powered, root-free theorem over finite sums, with an automatically checked certificate for every coefficient index in (III.44.4).*

V.85 Conclusion

Source 20, delivered Section 11.

The decisive local change is to keep two functions at the last Cauchy–Schwarz stage. That produces a bilinear dependence on their uniformity norms, with one additional L^2 factor. An exact second-order expansion then converts the analytic estimate into quadratic symmetric counting error. Torsion is neither ignored nor excluded: the kernel-image directions identify it, a sharp

index comparison quantifies it, and explicit collapsed-pair examples show why some dependence is necessary.

The three-term construction also supplies a boundary to this line of improvement. Over admissible growing-dimensional groups, even genuine sets can have count excess proportional to the square of a small balanced U^2 norm. Thus further progress must exploit extra structure, a sharper coefficient, a different norm, a restricted ambient family, or information unavailable in the present symmetric estimate.

V.86 Source audit and package scope

Source 20, delivered Appendix A.

The source inspection took place on 6 October 2026. A repository branch checkpoint observed during the work was

1bd57e91c76800c175c311fd1f37b2a10f07b9db.

The repository was active during inspection. The earlier research comparison is therefore identified by its actual inspected blob, not silently treated as the complete state of every later commit.

Source	Inspected identifier or range
Original paper TeX in the Ramsey paper directory	41085a0efde8c1932e86e80c791984f52841c22e. Paper Theorem 3.2, Corollary 3.3 and uniformity normalization were also checked against the original PDF, including rendered pages.
Existing local-refinements article	45b408f6d970a6d997973189663395670b5f3e6e. Introduction, cube/sharpness discussion, and the density-weighted progression section, especially source lines 1840–2025.
Sections01_03.lean	[write] Source 04’s delivered manuscript; the first write replaced <code>article.tex</code> at 14:34, after this inspection. f9fcd166af0ea5a8e8c674576c1e47f3705485cd, inspected at the checkpoint above, including definitions on lines 250–340.
Proofs01_03.lean	e72f7bd53a23d5880cbddbfefd0664d529e47c10, Fourier-identity infrastructure inspected near the start of the file.

The source manifest records these identifiers and the primary external references. The archive includes the article source and PDF, verification program and report, build instructions, and this manifest. It does not redistribute Gowers’s paper, the repository, or third-party font files. No repository write, external submission, or new formalization claim is part of the package.

V.87 Repository interface and verification

Source 21, delivered Section 10.

V.87.1 What was inspected

The source audit used the GitHub connection to inspect `VladimirReshetnikov/ProveIt` at commit `43306653afaaaaa0b58a7ca71e4067c245594867`. The inspected directory was `Combinatorics/Ramsey`. The principal comparison files were the Gowers transcription, the existing research compilation, the Section 3 statements and definitions, the cube upper-bound proof, and the proof-status ledger. The frozen commit is used in the source manifest so later edits of the moving branch cannot change the object being compared.

Source or declaration	Role in the comparison
Gowers, Lemma 3.8	Cube Cauchy–Schwarz, the standard analytic input
Gowers, Lemma 3.9	Triangle inequality; the projection consequence quantifies a related equality defect
Gowers, Lemma 3.10	The general cube upper bound strengthened in the real centered characteristic-five setting
<code>Sections01_03.lean</code> , <code>lemma_3_10</code>	The formal statement of the original cube upper bound
<code>Proofs03CubeUpper.lean</code> , <code>lemma_3_10_holds</code>	An existing exact companion declaration, as also recorded in the ledger
<code>Definitions.lean</code>	Conventions for <code>UniformOfDegree</code> , <code>gowersNorm</code> , and <code>cubeCount</code>
Earlier extremal manuscript, Section 3	The specifically unresolved characteristic-five coefficient question

The ledger’s status for `lemma_3_10` is `exact_companion`. This is a report of the inspected source and ledger, not a fresh build certification. The new inequalities in this article are written mathematical proofs; no new Lean theorem is claimed, and the repository was not modified by this task.

[write] **6 October 2026; sources 21–27.** *Checked at HEAD `2dfe46f44`.* `lemma_3_10` is at `Sections01_03.lean`, line 340, and its exact companion `lemma_3_10_holds` at `Proofs03CubeUpper.lean`, line 82; the ledger still records `lemma_3_10` as an exact companion. The seven repository files this source cites have the same blobs at HEAD as at its pin, so its comparison is current. No characteristic-five or cube-expansion declaration exists.

V.87.2 A concrete formalization route

The shortest route to formalizing the new material is to separate finite combinatorics from real optimization.

1. Prove the Fourier sum-of-squares identity on a finite abelian group and specialize to real centered functions. Establish the axis identity as a finite-sum equality.
2. Formalize the three opposite-pair families and the eight extra order-five tuples as disjoint weighted finite sets. This gives (III.52.2) without analytic asymptotics.
3. Prove the three six-vertex identities and the derivative identity for C_7 algebraically. The displayed profile inequalities then use finite-sum Cauchy–Schwarz and Hölder.

4. State the finite profile theorem (III.52.4) with explicit nonnegative real parameters. Keep this usable theorem independent of asymptotic notation.
5. Prove the scalar localization and absorption lemmas, and then the two-variable quadratic optimization. This isolates the absolute smallness threshold and remainder constant.
6. Treat the five-point polynomial as a finite certificate, followed by a separate real analytic optimization theorem. A formal proof of its exact coefficients is substantially easier than a formal proof of the analytic maximizing branches.

The general exponent-five group requires a finite abelian-group interface or a vector-space model over $\mathbb{F}_5$; a statement only on $\mathbf{ZMod}\ N$ does not automatically represent $\mathbb{F}_5^n$. Normalization also matters. In this article the cube average is N^{-4} times the count and the centered uniformity budget is u^8 . Any implementation should verify these factors at the interface, rather than transfer theorem names alone.

V.87.3 Computational checks supplied with the article

The archive contains two standard-library Python verifiers. Their role is to test finite algebra and catch transcription errors; the proofs above establish the general results.

- `verify_order5.py` enumerates all 625 Fourier cube assignments on $\mathbb{F}_5$, obtaining 34 monomials. It compares physical-space cube averages, the full Fourier formula, the six-row table, and the expanded phase polynomial. It checks the rational series coefficients and $20/9$ exactly.
- `verify_profiles.py` checks the three sextic support identities and the septic derivative by exact monomial dictionaries on $\mathbb{F}_5$. It also checks the finite profile and energy inequalities on deterministic examples in $\mathbb{F}_5$ and $\mathbb{F}_5^2$, including direct physical-space cube averages.

The supplied run records report all checks passing. The largest normalized floating-point discrepancy in the two independent runs was below $2 \cdot 10^{-13}$. This small numerical discrepancy is consistent with rounding; the exact dictionary and rational checks use integer or rational arithmetic. The general asymptotic claims are not inferred from finite sampling.

V.88 Further research questions

Source 21, delivered Section 11.

The questions below distinguish a natural next calculation from a larger structural problem. None is asserted to be a previously published open problem without a separate source establishing that status.

Question V.88.1 (The universal eighth-order coefficient). *For fixed $0 < \delta < 1$, let*

$$M_\delta^{(5)}(u) = \sup_{\substack{G \neq 0 \text{ finite abelian, } 5G=0 \\ W: G \rightarrow [0,1], \mathbb{E}W=\delta \\ \|W-\delta\|_{U^3} \leq u}} Q_3(W).$$

Does the limit

$$\lim_{u \downarrow 0} \frac{M_\delta^{(5)}(u) - \delta^8 - 6\sqrt{2}\delta^4 u^4 - C_5 \delta^2 u^6}{u^8}$$

exist, and is it 20/9? Theorem F.26.1 gives a finite upper limsup, while Corollary III.55.2 gives lower liminf at least 20/9. The missing step is a sharp treatment of frequencies outside the dominant five-point quotient.

Question V.88.2 (Fixed higher dimensions). For $G = \mathbb{F}_5^n$ with fixed $n \geq 2$, is the eighth-order coefficient always 20/9? Can an extremizing family be proved to depend on one linear functional for all sufficiently small norms? The projection estimate (III.57.1) supplies approximation, but not exact attainment after norm rescaling.

Question V.88.3 (A quantitative eighth-order stability theorem). If the objective is within $O(u^8)$ of the displayed sharp sixth-order expansion, what is the best bound for the Fourier mass outside $\{\pm\xi, \pm 2\xi\}$? Determine whether a U^2 error of order u^2 is the correct uniform scale and identify the additional phase and amplitude losses.

Question V.88.4 (Further terms on the five-point group). The analytic maximizing branches in Theorem III.55.1 permit recursive calculation of higher coefficients. Determine the exact u^{10} and u^{12} terms and the first nonzero displacement of the phases. At order u^{12} , the phase perturbation can affect the value, so a computation confined to the limiting phases is insufficient.

Question V.88.5 (Exact finite-norm maximizers). For fixed $0 < \delta < 1$, how far can the analytic maximizing branches be continued before a weight reaches 0 or 1 or a different branch takes over? The six-row formula reduces the question on $\mathbb{F}_5$ to a compact, explicit constrained problem. Certified interval arithmetic or elimination could locate its transition values.

Question V.88.6 (Other torsion classes). Determine the comparable near-constant expansion on exponent-three groups, where the centered low-degree cancellations and pure-pair norm differ. For mixed odd torsion, classify the secondary behavior according to the order of the largest Fourier character, taking care that competitors of several character orders may coexist.

Question V.88.7 (Higher-dimensional additive cubes). For $d \geq 4$, determine the sharp leading real-centered coefficient in $Q_d(\delta + f)$ at fixed $U^d(f)$, and classify the first torsion-dependent correction. The three-cube Fourier sum-of-squares representation is unusually simple; higher dimensions require a larger incidence-kernel analysis.

Question V.88.8 (Optimal profile inequalities). The estimates $|C_6| \leq 28H$ and $H \leq 3\sqrt{DE(\lambda)}$ are adequate for the sharp sixth-order theorem, but their constants and simultaneous equality cases have not been optimized here. Determine which profiles can actually approach equality in all the needed estimates.

Question V.88.9 (Quantitative conversion from weights to sets). The probabilistic argument allows the ambient dimension to be made arbitrarily large. What is the least dimension needed to approximate the model through a specified order, or to realize a specified density with an explicit cube-count error? Deterministic rounding would provide a useful constructive complement.

Question V.88.10 (Machine-checked finite certificates). Can the 34-monomial five-point certificate and the sextic support identities be checked by a small trusted Lean evaluator, followed by a proof of the scalar profile theorem? This is a concrete entry point that does not require formalizing the entire analytic maximization argument at once.

Question V.88.11 (Where local gains affect larger arguments). *Which finite-field density or counting arguments genuinely use this sharper cube count with its characteristic-five structure? Identify an explicit downstream theorem before claiming a quantitative improvement. The absence of long proper progressions in exponent-five groups prevents a direct extrapolation to integer Szemerédi bounds.*

The most focused next target is the first question: the proofs already isolate a finite eighth-order gap, and the model supplies the exact candidate 20/9. A resolution would require a sharper geometric analysis of the off-line Fourier tail, rather than another optimization of the two-harmonic phases.

V.89 Finite incidence certificates and reproducibility

Source 21, delivered Appendix A.

V.89.1 Centered support census

Associate to each cube vertex ω the column $(1, \omega_1, \omega_2, \omega_3)^\top$ over $\mathbb{F}_5$. The expectation of a product of centered functions on a support is the sum of Fourier monomials whose frequencies lie in the kernel of the corresponding column matrix. A zero forced coordinate makes the monomial vanish. The relevant small census is

Support size	Number	Classification relevant to the expansion
1	8	Independent; centered expectation zero
2	28	Independent; centered expectation zero
3	56	Independent; centered expectation zero
4	70	12 rank-three parallelograms; 58 rank four
5	56	8 full-support circuits; 48 with a forced zero
6	28	Omitted-pair distances 1, 2, 3: 12, 12, 4
7	8	One omitted vertex; the derivative formula applies
8	1	The full U^3 power

This census explains why the first three errors vanish and why the fourth and fifth terms have coefficients 12 and 8.

V.89.2 The full five-point polynomial certificate

Order the vertices as 000, 001, 010, 011, 100, 101, 110, 111. All 625 solutions of the four Fourier constraints on the full cube can be parameterized by $(s, t, v, w) \in \mathbb{F}_5^4$ as

$$(s + t + v + 2w, -s - t - w, -s - v - w, s, -t - v - w, t, v, w).$$

The four equations can be checked by substitution, and the last four chosen parameters are recovered directly from the labels at 011, 101, 110, 111, proving that the parameterization is bijective. Counting the multiplicities of the monomials $a_0^{e_0} \cdots a_4^{e_4}$ gives the accompanying CSV certificate. The multiplicities sum to 625, and every exponent vector has total degree eight.

This is independent of the six-slot sum-of-squares calculation: one certificate solves the physical cube’s orthogonality equations, while the other groups Fourier parallelograms into squared terms. Their agreement supplies a useful finite algebra check.

V.89.3 How to reproduce the supplied artifacts

From the extracted archive directory, run

```
python3 code/verify_order5.py
python3 code/verify_profiles.py --output data/profile_validation_results.json
make pdf
```

The verifiers need only Python’s standard library. The optional figure-generation program uses `matplotlib` with high-precision arithmetic from Python’s `decimal` module; the prebuilt PDF figure is included, so plotting packages are not needed to compile the article.

V.90 Source manifest and status of claims

Source 21, delivered Appendix B.

The machine-readable source manifest records the repository commit, paths, and Git blob hashes of the inspected files, together with the primary literature URLs. The earlier project report was read as a separate manuscript, not assumed to be part of the inspected repository snapshot.

Claim category	Status in this article
Gowers–Cauchy–Schwarz and norm property	Standard results of Gowers; used with their normalization stated explicitly
Quartic gap and earlier characteristic-five lower example	Prior project results, reproved or recovered from exact formulas
Uniform coefficient C_5 with $O(u^8)$	Proved in Theorem F.26.1, with matching examples
Phase envelope and necessary rigidity	Proved in Theorems III.53.2 and III.54.1
Five-point coefficient 20/9	Proved for the fixed group in Theorem III.55.1
Universal eighth-order coefficient	Undetermined; 20/9 is a lower bound and an explicit candidate
New Lean certification or new global Szemerédi bound	Not claimed
Bibliographic priority	Not established by the source search

The proof development and an independent mathematical audit agreed on the core inequalities and both asymptotic optimizations. The article incorporates the audit’s corrections concerning the zero norm case, the sign of a phase cosine, and consistent phase notation. The included review note records that mathematical review as an AI-assisted check; it is not external peer review.

V.91 Verification, attribution, and repository integration

Source 22, delivered Section 11.

V.91.1 Exact computational checks

The package contains four small Python programs, using only the standard library. They use exact integer, rational, and finite-field arithmetic. The results are recorded as JSON so that the numerical claims can be reproduced independently of the typeset manuscript.

Program	Verification performed
<code>verify_certificates.py</code>	Exhaustive moment–degree duality for all ternary slices in dimensions one through three over $\mathbb{F}_5$ and $\mathbb{F}_7$; affine slice counts; independent sparse Laurent convolution versus the joint-moment recurrence; the explicit $d = 2$ polynomial and a $d = 3$ binomial-basis polynomial; exact coefficient comparison tables.
<code>verify_two_cubes.py</code>	Direct enumeration of all ternary slices and normalized translations in the listed $k = 1, 2$ cases. It matches the full squarefree moment vectors and checks the stated finite interval for $\delta_{k,1}$, without using the stabilizer argument.
<code>verify_polynomial_patterns.py</code>	20 400 norm evaluations and all 74 norm-value pairs over $\mathbb{F}_5, \mathbb{F}_7$ certify the origin examples; 512 Vandermonde-budget cases include degrees exceeding characteristic; 3 476 quadratic maps and 2 051 540 center/direction evaluations check the quadratic classification; 444 extra-direction fibers at 72 centers check the cubic-core formula using 119 592 core directions.
<code>verify_norm_obstruction.py</code>	Irreducibility of the field moduli, all eligible ratios for $(q, p) = (3, 3), (5, 5), (9, 3)$, and an independent full start/direction enumeration in $\mathbb{F}_{27}$. The ratio counts are 6, 20, 18; the corresponding full- p progression counts are 156, 62 480, 13 104.

The factorized origin check for $N_3 + N_5$ is exact: every norm value has its verified fiber size, and every pair of norm values is tested against the two reflection equations. It therefore certifies a count over all q^8 directions without pretending to have explicitly iterated that many vectors. Likewise the value 7 810 000 is a proved lower bound for the global symmetric count, not a numerical estimate or an exact global count.

The two-cube checks found no extra degeneracy outside the listed hyperplanes in their tested cases. The theorem uses an explicit remainder and does not promote that observation to an all-prime statement. All general claims are proved in the body of the article. Independent internal review checked the degree filtration, the class quotients, the finite summation by parts, the zero-set rigidity, and the distinction between prime-field and extension-field quadratic functions.

V.91.2 How this fits the existing formalization

At the inspected repository snapshot, the arrangement definitions are in `Combinatorics/Ramsey/Lean/GowersSzemerédi/Sections14_15.lean`. They include `GeneralArrangement`, `arrangementMoment`, `GeneralArrangement.IsDegenerate`, and `degenerateGeneralArrangementCount`. The existing file `Proofs15DegenerateArrangements.lean` contains `lemma_15_4_holds` for the original-size estimate. The present arrangement results use the same full moment definition of degeneracy, not a weaker collision-only predicate.

A practical formalization order is as follows.

1. Define parity-normalized certificate degree by the finite Boolean interpolation coefficients, then prove Lemma II.82.2 as an identity of finite sums.
2. Prove the diagonal and paired-diagonal rank statements of Lemma II.82.3; an injective choice of the pivot variables gives the finite counting bound directly.
3. Formalize affine ternary slice classification and the affine-value equivalence relation. First use U_1 ; the slope-balanced constant-term refinement can then be added without changing the rank proof.
4. Prove the cumulative weighted-sum lemma once in a general ordered semiring or real setting. It is a reusable guard against invalid shell subtraction in other counting arguments.
5. Add the refined arrangement theorem under its explicit prime condition $p > 4d$, and retain the separate $d = 1$ theorem under $p > \max\{3, 2k + 1\}$.
6. Add polynomial equality patterns as a separate definition. A symmetric four-term pattern and a monochromatic four-term pattern have different predicates and different degree budgets. Their identifiers should make that distinction explicit.

This package contains mathematical proofs and executable finite checks, but no new Lean proof file. It has not rebuilt the repository or altered its proof-status catalogue. In particular, this document should not be used to mark a theorem as kernel-checked.

[write] **6 October 2026; sources 21–27.** *Checked at HEAD 2dfe46f44.* Every name above exists (`Sections14_15.lean`: `GeneralArrangement`, line 192; `arrangementMoment`, line 346; `GeneralArrangement.IsDegenerate`, line 354; `degenerateGeneralArrangementCount`, line 364), and `lemma_15_4_holds` is at `Proofs15DegenerateArrangements.lean`, line 1131. The suggested separate directory was declined: the manuscript is printed in this report (Section F.4).

V.91.3 Suggested placement and reproducibility

The standalone source `gowers_degree_stratification.tex` has no external bibliography or figure dependency. It can be placed in a new subdirectory of `Combinatorics/Ramsey/Research/GowersSzemerédi/` along with the accompanying `code/` and `data/` directories. The supplied PDF is compiled from that source. The README records the build commands and a concise statement-level comparison with the earlier reports. The source ledger records the repository revision, source paths, and checksums used to ground the work.

The intended integration is additive: a new research article and its verification artifacts. The theorem numbering in Gowers’ paper is not renumbered, and the existing formal statements are

not silently replaced. Each quantitative substitution should cite the new theorem together with its stronger characteristic hypothesis.

V.92 Further research questions and proposed next steps

Source 22, delivered Section 12.

The following questions distinguish exact problems exposed by this article from broader research goals. None is asserted to be a recognized major open problem solely because it is listed here.

Research question V.92.1 (The actual second degeneracy coefficient). For fixed k and $d \geq 2$, determine the coefficient of p^{-2} in $\delta_{k,d}(p)$. The present theorem reduces the dominant nonparity search to slope-balanced affine Boolean slices. A solution must still identify coincident codimension-two loci and their intersections with the pure-parity hyperplanes. Counting certificate classes alone cannot give the answer.

Research question V.92.2 (Exactness in the two-cube case). Is $\delta_{k,1}(p) = L_k(p)$ for every sufficiently large prime p , or do some nonaffine translation loci contain points outside $t_i = 0, \pm 1$ and $t_i \pm t_j = 0$? The finite checks for $k = 1, 2$ suggest a concrete starting point. Either a containment theorem or a single exact counterexample would sharpen the structural picture.

Research question V.92.3 (Degree-two Boolean certificates). Count ternary-valued Boolean polynomials of degree at most two together with their necessary moment balances, up to affine value changes. The first target is an explicit cumulative bound much smaller than $A_{2d(1+k+\binom{k}{2})}$. A successful classification should feed directly into the p^{-3} term of the finite bound.

Research question V.92.4 (A feasible-class filtration in every degree). The current higher-degree profile counts many classes whose lowest moment constraints are inconsistent. Can one build cumulative bounds after imposing all necessary constraints, while retaining a valid summation-by-parts comparison? The exact degree-one endpoint shows why this is promising and also why differences of separate upper bounds must be treated carefully.

Research question V.92.5 (The balanced constant-term polynomials). Find a compact generating function in both k and d for $J_{k,d}$. The leading coefficient is an alternating-cycle count; the lower coefficients should admit a description using a finite collection of decorated incidence graphs. Such a description could provide efficient exact values for the eight-arrangement application.

Research question V.92.6 (Small characteristics). Extend the balanced count to primes $p \leq 4d$ by replacing integer constant terms with zero coefficients in a finite group algebra. Which additional modular balance types change the leading or second degeneracy coefficient? The Boolean degree/rank argument itself remains available, but the pattern-class and no-aliasing counts must be redone.

Research question V.92.7 (A fully propagated quantitative bound). Carry the improved exceptional-set function through the entire chain from respected arrangements to multilinear structure, phase removal, density increments, and iteration. The correct output is a checked parameter dependency graph with a closed final estimate. A local improvement alone does not establish a better global Szemerédi bound.

Research question V.92.8 (Global quintic avoidance beyond the cubic core). Determine the maximum dimension of a scalar polynomial of degree at most five, or six, avoiding symmetric

four-term progressions over fields of characteristic greater than four. The current range is four to eight. The exact pointwise theorem rules out improving eight by an argument that guarantees a nonzero direction at each independently prescribed center. The no-additive-extension theorem rules out the entire family $x_0 + N_3(u) + R(w)$ with added variables. Mixed dependence between the old and new coordinates is therefore a natural construction target.

Research question V.92.9 (Simultaneous sharpness for a general equality pattern). For which node partitions and degree caps is the minimum $q^{\max(n-B,0)}$ attained by one polynomial at every center? The low-degree monochromatic and symmetric patterns give positive examples. The high-degree norm example shows that pointwise sharpness alone supplies no answer.

Research question V.92.10 (Classifying general pointwise extremizers). Minimum support determines the geometry of the successful direction set, but not all polynomials producing it. Classify the maps that attain the pointwise minimum for a given evaluation matrix. The complete quadratic normal form is a first instance. A useful classification should account for centers at which the subspace changes.

Research question V.92.11 (Surjectivity beyond complete norm cores). Characterize polynomial cores whose reflection map is onto at every center. Such a characterization would identify a broad family for which independent additive extensions cannot improve avoidance. It is also natural to ask for lower bounds on every reflection fiber, not only surjectivity.

Research question V.92.12 (Higher-order extension defect formulas). The cubic-core formula expresses the exact global count in terms of the support of a third difference. For five or more odd norm layers, can one obtain an equally useful expression in a hierarchy of odd differences? The triangular norm construction supplies a recursion, but higher-block choices affect the targets of lower blocks.

Research question V.92.13 (Stability of the zero-defect classification). If the third difference of R vanishes on almost every pair, is R close to a prime-field quadratic function with a quantitatively sharp error bound? The exact cubic-core formula converts this directly into stability for the number of symmetric progressions in independent extensions. Connections with classical low-degree testing should be used explicitly rather than rediscovered without attribution.

Research question V.92.14 (Norm lines outside characteristic-divisible degrees). For general t , count ratios θ for which $N_t(\theta + j)$ is constant on $\mathbb{F}_p$, or merely satisfies the required reflection equalities. The exact degree- p answer uses a Frobenius orbit on an Artin-Schreier translate. Other degrees may require a more general description of the conjugate-root polynomial.

Research question V.92.15 (An initial Lean milestone). Formalize the exact formulas $J_{k,2}$ and $\delta_{1,1}(p)$, together with the abstract weighted cumulative-sum lemma and the evaluation-matrix minimum. These results have small interfaces and provide useful reusable components before the full degree-stratified arrangement theorem is attempted.

V.93 Verification record and repository integration

Source 23, delivered Section 9.

V.93.1 What the accompanying checks establish

The package includes three independent exact-arithmetic programs. They use Python’s standard library and make no network requests. Their purpose is to detect finite combinatorial, normalization, and transcription errors in the formulas used by the proofs. They do not replace the arguments establishing asymptotic uniformity or arbitrary-dimension inequalities.

Table V.93.1: Exact verification performed for this article.

Program	Verified content
<code>verify_prescribed_partitions.py</code>	1770 interval conductors; 20181 nonzero mixed-radix residues; 3122 rational phase vectors for the torus bound; four complete partitions; 5000 exact residue-packing formulas; 196 exceptional-mass gap calculations; and the prime obstruction with modulus 29269.
<code>check_fifth.py</code>	The entire centered cube polynomial, all six entries of the five-point Fourier-slot table, the exact norm and weighted-count polynomials, and the 12, 12, 4 sextic support multiplicities. Coefficients are integers in a formal phase ring with $z^{10} = -1$.
<code>verify_section16_gluings.py</code>	177103 localization instances, 88688 contained short progressions, 1099140 residue-transport instances, and 2361 sampling profiles and sample sizes, including the exact add-one identity.

The phase program verifies its full partitions by checking coverage, disjointness, allowed lengths and phase diameters, not just the chosen denominator. For the prime obstruction it enumerates every possible step allowed by the minimum cell length. In the cube program, the special phase is represented algebraically; floating-point cancellation is not used to certify the polynomial identities. The transport program checks both orientations of the integer step and the cardinality of every residue chain.

The included figure is generated from proved upper and lower formulas. Its shaded band bounds an unknown exact threshold. It is not the result of optimizing over all phase vectors numerically.

V.93.2 Suggested integration into the research collection

This article can be placed in a new research directory, for example

Combinatorics/Ramsey/Research/GowersSzemerédi/
sharp-thresholds-and-extrema/.

Keep the article as a separate unit before merging it into the large collection: its central contribution is the completion of specific sharpness questions, and this is easiest to review when the earlier sufficient bounds remain visibly attributed.

In the research-question ledger, the following updates are justified by the written proofs:

1. Mark the prescribed-length conductor-loss question as resolved for simultaneous affine phases with arc diameters below $1/2$, in the eventual universal-threshold sense. Record both the finite sandwich and the leading constant one. Keep separate any question about a particular interval length, an exact number of cells, or nonlinear phases.

2. Mark the characteristic-five secondary cube coefficient as determined: it is $2^{3/4}/3$ with an absolute $O(u^8)$ error. The result applies to real centered functions, gives the weighted supremum on every fixed nontrivial exponent-five group, and has indicator sharpness in growing dimension. Keep the tertiary coefficient open.
3. Record the common-base construction and full-width affine transport as proved conditional lemmas for the proper modular-box interface. Keep the final Section 16 control-function comparison and the ordinary-integer source-fidelity problem open.

The repository’s original paper transcription is an object of ongoing mathematical repair. The result here should be entered as an independent theorem and linked to the relevant numbered lemma; it should not silently replace the source’s statement or erase the history of a correction.

V.93.3 A feasible formalization sequence

The partition results have a particularly direct finite formalization route. Define the semigroup of permitted cardinalities, establish its conductor property, and prove the interval conductor formula. The mixed-radix obstruction needs only divisibility, centered modular distance, and counting residue classes. It would already certify the new sharpness conclusion without formalizing a measure-theoretic torus packing argument. The slightly weaker rectangular upper bound is entirely finite and has the same sharp asymptotic; the stronger finite packing constant can be added later.

For the cube theorem, a useful interface consists of the normalized Fourier identity (F3), the exact homogeneous expansion (F4), the concentration penalty (F8), and the incomplete-cube bounds (F9). The subsequent optimization is a real inequality with absolute small-parameter hypotheses. The five-point lower model can be checked by a finite algebraic certificate, but a proof assistant still needs an explicit bridge from that certificate to the Fourier definitions.

For geometric gluing, the natural order is Lemma II.103.3, Lemma II.103.4, Theorem II.103.6, and finally Theorem II.103.7. The existing modules `Proofs16RelativeSteps.lean`, `Proofs16CommonBase.lean`, `Proofs16Interpolation.lean`, and `Proofs16SampleSelection.lean` provide nearby interfaces [110]. The proofs must preserve the orientation of a negative integer lift and the exact cardinalities of proper progressions. No claim that these proposed declarations have been kernel-checked is made in this package.

[write] **6 October 2026; sources 21–27.** *Checked at HEAD 2dfe46f44.* The four modules named above exist; `lemma_16_10` (`Section16.lean`, line 851) is open, and the formal results compared with Section II.103 are in the notes there. The research-question updates proposed in this section are made in Section V.2 (items I.12, III.23 and II.40). The suggested separate directory was declined: the manuscript is printed in this report.

V.94 Further research questions and topics

Source 23, delivered Section 10.

The first two questions have the shortest path from the present proofs to additional quantitative results. The later questions address structural generalization and the unresolved global dependency chain.

1. **Determine the second asymptotic term of the prescribed-length threshold.** We now know $T_{\{L, L+1\}}(\epsilon) = L^{q+2} / \prod_i \epsilon_i (1 + O_q(L^{-1}))$. The lower and upper bounds differ at

the relative $1/L$ scale. Is there a universal next coefficient, or do rounding and arithmetic properties of the tolerances force oscillation? Begin with $q = 1$ and a fixed rational tolerance, where every relevant step can be inspected exactly on small intervals.

2. **Determine the tertiary cube coefficient in characteristic five.** The sharp secondary coefficient leaves an $O(u^8)$ error. At this order, additional Fourier modes can compete with the two selected harmonics. The exact five-point polynomial gives an explicit test family; a rigorous global optimum would need a sharper treatment of the tail terms than the union bound in (F9). A finite-dimensional variational reduction would be particularly useful.
3. **Obtain a quantitative stability theorem for the phase obstruction.** The mixed-radix construction realizes the sharp asymptotic threshold. Must nearly worst phase vectors resemble a sequence of nested rational quotients after a suitable change of coordinates on the phase torus? Specify a metric and a notion of near-extremality before attempting the classification; a threshold statement alone does not force such a structure.
4. **Find exact thresholds when the range of lengths is a fixed fraction of the minimum.** When $w = o(L)$ the leading constant is known. For $w \sim \theta L$, $0 < \theta \leq 1$, the conductor is explicit but the present upper and lower phase factors need not have the same leading constant. One can first optimize the upper construction over shorter prefixes of the permitted interval. A matching lower construction may require more information than the smallest allowed length.
5. **Classify optimal exceptional-set profiles.** Formula (I.39.1) is exact for the rigid examples, and (I.39.2) determines the joint limit $L \rightarrow \infty$, $\eta \rightarrow 0$ for one exact length. What is the leading constant for fixed $\eta \in (0, 1)$? How does the answer depend on a general numerical semigroup of allowed lengths? The gap function g_S supplies an explicit family of lower bounds.
6. **Extend the obstruction to polynomial phases with prescribed lengths.** A degree- d phase on a progression is controlled by several finite differences. Seek a higher-degree analogue of the mixed-radix construction that forces progression steps into a useful divisibility class while retaining the tiling obstruction. The goal is a matching necessity theorem for a carefully stated polynomial-partition bound, rather than only an upper recurrence estimate.
7. **Quantify the characteristic-five phase rigidity.** The proof forces harmonic alignment and $\cos(5\theta) \rightarrow 0$ at optimal secondary order. Turn the deficit below $2^{3/4}/3$ into explicit bounds for $|b/(\kappa\delta^{-1}u^2) - 1|$, the relative phase error, and the fourth-power Fourier tail. A bound in L^2 would require additional information and should not be inferred from these tail estimates.
8. **Analyze mixed torsion groups.** The norm penalty caused by order five is tied to the order of the dominant character. In an odd-order group supporting both order-five and larger-order characters, determine which spectral regime maximizes the local cube count under additional constraints. The unrestricted supremum and a supremum conditioned on the dominant character order are different problems.
9. **Improve common slice refinement through joint structure.** The exponent in Theorem II.103.6 is a product over all sampled slices. A structural theorem for the vector of slice values might avoid some of this repeated loss. The target should give one common partition directly, with explicit exceptional mass and list size, so that interpolation can use it without another refinement.

10. **Optimize localization when the child width is large.** The factor $1/3$ is sharp for the width-independent cap used here. The transport proof actually needs $tw \leq |J|$, where $t(u-1) \leq b-1$. When a large lower bound for w is already known, the worst width-two obstruction disappears. A width-dependent localization rule could improve the constant and perhaps the finite thresholds used in Section 16.
11. **Close the Section 16 parameter recurrence.** Substitute the actual affine-class count, anchor count, exceptional budgets, and slice profiles into the gluing theorem. Determine control functions that are closed under every step of the induction. The source’s displayed final comparison is not justified by the local construction alone. A useful deliverable would be a checked recurrence with explicit initial values and a proof of its solution, followed by a complete dependency audit.
12. **Formalize the sharp lower bounds first.** The numerical semigroup and mixed-radix arguments are finite, relatively short, and independent of the difficult structural steps in the original proof. They provide a realistic first kernel-checked addition. Then formalize the exact cube identities and the transport lemma. Maintaining a clear distinction between source statements, repaired statements, and new refinements will make later integration substantially easier.

V.95 Build instructions and package contents

Source 23, delivered Appendix B.

The ZIP contains one complete LaTeX source, its compiled PDF, the figure files, three exact verification programs, their recorded outputs, a figure-generation script, a source manifest and a Makefile. The article does not require access to the original papers or previous reports in order to compile. All mathematical sections are in the single file `gowers_sharp_thresholds.tex`.

With a standard TeX Live installation, run

```
pdflatex -interaction=nonstopmode -halt-on-error gowers_sharp_thresholds.tex
pdflatex -interaction=nonstopmode -halt-on-error gowers_sharp_thresholds.tex
pdflatex -interaction=nonstopmode -halt-on-error gowers_sharp_thresholds.tex
```

Alternatively, `makepdf` runs these three passes. The exact checks are reproduced by `makecheck`; they need Python 3 and no external packages. Regenerating the optional figure requires NumPy and Matplotlib, but the generated PDF figure is already included. The source manifest records the inspected repository revision and source URLs. No Lean source or compiled proof certificate is included.

V.96 Examples, boundary cases, and reproducible checks

Source 24, delivered Section 10.

V.96.1 Three structural examples

Example V.96.1 (An affine map has no vertical obstruction). If $\phi(x) = T(x) + a$ on B , where $T : G \rightarrow H$ is a group homomorphism, then $F_k = \{0\}$ for every k . Equal domain sums produce

equal output sums because both sides contain the same number of copies of a . The extraction theorem therefore returns all of B , without paying the generic factor kC^{2k+1} .

Example V.96.2 (A small nontrivial fibre with an active avoidance step). Let $p = 101$, $B = \{0, 1, 2\}$, $k = 2$, and $\phi(0) = \phi(1) = 0$, $\phi(2) = 1$. The only conflicting two-term domain sum is $0 + 2 = 1 + 1$, so

$$F_2 = \{0, 1, -1\}, \quad Q_2 = 3.$$

The avoidance length in (II.79.5) is $L = \lfloor 99/(2 \cdot 2) \rfloor = 24$. The choice $d = 2$ works: the multiples $2j$ with $1 \leq j \leq 48$ are neither 1 nor -1 modulo 101, and symmetry handles negative j . One colour fibre contains $\{0, 1\}$. This is an order-two Freiman restriction, while the entire map is not. The example exercises the nonconstant-output, nonzero-length part of the proof rather than merely the $L = 0$ pigeonhole case.

Example V.96.3 (Different objectives select different cores). Theorem II.74.1 is appropriate when the original threshold $\delta^2 m/2$ and the 90% conclusion must be preserved. For an additive application, the one-sample choice $\beta = p^2/16$ trades a smaller common-neighbour count for a larger core. The popular-mass cancellation then makes that trade profitable. There is no universal rule that more samples improve the final structural estimate.

V.96.2 Finite checks supplied with the article

The companion `verify.py` uses only Python’s standard library. All inequalities tested are evaluated with integers or exact rational numbers. The recorded run completed the following checks:

Check family	Cases
Cubic supporting-line inequalities on a rational grid	4,040
All incidence matrices with four indices and three ambient points	4,096
Weighted moment identities with deterministic random inputs	100
All nonempty symmetric four-vertex graphs, allowing loops	1,023
All nonempty subsets of $\mathbb{Z}/q\mathbb{Z}$, $q \in \{3, 5, 7, 11\}$	2,212
Graph maps for vertical-fibre and extraction checks	3,216

Table V.96.1: Exact finite regression checks. All passed. Among the graph-map checks, 60 exercised $Q_k > 1$ and $L > 0$ simultaneously.

For the incidence matrices, the program enumerates every ordered three-sample, verifies the exact expectation identity, and checks that the best score is at least $49u^3n^2/54$ with u equal to the actual mean codegree. For each small graph it verifies the one-sample score, the pruning inequality, and the common-intermediate walk bound. For each small additive set it checks both popular-mass estimates and all three rounded conclusions of Corollary II.77.2.

The fibre checks verify the injection into $(k + 1)\Gamma - k\Gamma$, the cardinality inequality, the existence of an allowed scalar, and the Freiman property of the constructed colour fibre. All full-domain maps $\mathbb{F}_p \rightarrow \mathbb{F}_p$ for $p \in \{2, 3, 5\}$ are tested at order two, together with 60 three-point-domain maps over $\mathbb{F}_{101}$ and $\mathbb{F}_{503}$. The latter are generated by a fixed random seed; the checks themselves remain exact.

These experiments can catch sign errors, threshold errors, and incorrect constants. They do not prove the theorems for arbitrary sets or all orders. In particular, the finite-field-vector hashing theorem is proved above but is not part of the supplied finite regression suite. The JSON report states the scope explicitly.

V.96.3 Constructive content and its limits

The core-selection proofs are finite algorithms. One can compute codegrees, score every allowed sample, choose a maximum-score sample, and prune its bad-degree vertices. In the one-sample argument there are only m candidate samples, once the neighbourhood table is known. The four-walk tuple injection proves the difference bound; an implementation need not enumerate D^4 to select the set.

The extraction step is also constructive after F_k is known: compute L , search the nonzero scalars for an allowed d , and take a largest colour fibre. Computing the exact defect fibre can itself be expensive as k grows. No claim of an efficient general-purpose algorithm for arbitrarily large k is implicit in these existence proofs.

V.97 Integration with the inspected Lean development

Source 24, delivered Section 11.

V.97.1 The inspected snapshot

Repository observations refer to commit

2e0553b0c8a6c9c93221156cc83e95d68ee8bb07.

The relevant directory is

Combinatorics/Ramsey/Lean/GowersSzemerédi.

The statement module `Sections06_07.lean` contains `HasDenseIntersectionCore`, `lemma_7_4`, `lemma_7_5`, and `corollary_7_6`, with the source bounds listed in Table F.29.1. Its `proposition_7_3` is stated for finite sets in an integer-vector group. The module explicitly records the primality and density hypotheses needed at the corresponding interfaces [113].

The proof file `Proofs07DRC.lean` uses `DRCIndex := Fin 5`. It defines a common-sample core, a bad-pair count, and a score equal to the square of the core cardinality minus ten times the bad-pair count. Its finite-counting lemmas express sums over samples as fifth powers of codegrees. These are the exact identities that should be parameterized, not discarded.

The file `Proofs07BalogSzemerédi.lean` uses popular *sums* instead of popular differences, while citing `lemma_7_4_holds` as its selection input. Its early auxiliary lemmas include representation-count bounds and an energy identity. The use of sums is compatible with the argument here, but requires a sign-adjusted walk identity; it should not be changed silently.

[write] **6 October 2026; sources 21–27.** *Checked at HEAD 2dfe46f44.* `HasDenseIntersectionCore`, `proposition_7_3`, `lemma_7_4`, `lemma_7_5` and `corollary_7_6` are at `Sections06_07.lean`, lines 156, 135, 167, 183 and 197; `DRCIndex` at `Proofs07DRC.lean`, line 21, and `lemma_7_4_holds` at line 536. All four files have the same blobs at HEAD as at this source’s pin. The form of `proposition_7_3` in every abelian group is formal as `proposition_7_3_group_holds` (`Proofs07BalogSzemerédi.lean`, line 854); see the note in Section F.29.

V.97.2 A minimal mathematical interface for the refactor

A parameterized DRC layer should expose the identity

$$\sum_{v \in V^t} |K_v|^2 = \sum_{i,j} |A_i \cap A_j|^t, \quad (\text{V.97.1})$$

together with its bad-pair analogue. Both identities are integer counting statements before normalization. A separate real-valued lemma can then turn a positive score average into an existential core. This separation avoids interleaving casts, sampling combinatorics, and nonlinear real inequalities.

For the cubic specialization, the main new arithmetic obligation is (II.74.4). A rational formulation can clear denominators and use the factorization

$$s^3 - 3b^2s + 2b^3 = (s - b)^2(s + 2b), \quad b = 7u/6.$$

The negative branch is handled by monotonicity on $[0, u/2]$ and the endpoint value $u^3/108$. The rational constant $49/54$ is preferable to the algebraic optimum for a first implementation.

A graph-certificate layer should take as inputs the exact values $|D|$, $|K|$, the bad-degree threshold, and the common-neighbour threshold. It should return the injection into four labelled edge values and the relative difference bound. Only the final numerical corollary should substitute $\tau = \kappa/2$ or replace p by a lower bound.

V.97.3 Retaining the repository's popular-sum convention

For the sum graph, an edge between a and b is labelled by $a + b$. A walk a, b, v, c, a' then gives

$$(a + b) - (b + v) + (v + c) - (c + a') = a - a'.$$

The four labels still recover the intermediate vertices from fixed endpoints, so the same fibre injection and cardinality estimates apply. The energy identity and the estimate $\tau m|D| \leq pm^2$ use popular sum representations instead. This provides a direct route to the existing module without a wholesale rewrite of its representation-count infrastructure.

V.97.4 The fibre layer and its formal obligations

The proposed fibre lemma should be stated for a graph in a product of abelian groups. Its proof requires only finite-set membership in an iterated sumset and injectivity of

$$(y, (x, \phi(x))) \longmapsto (x, \phi(x) + y).$$

The nonempty-graph hypothesis must be retained when cancelling $|\Gamma|$. The growth estimate should then be imported in the exact form (II.78.3), with exponent $(k + 1) + k$.

The prime-cyclic avoidance layer should separately prove that $1 \leq j \leq kL < p$ implies invertibility in $\mathbb{F}_p$, count forbidden nonzero scalars, and establish the integer-division inequality $k(Q - 1)(L + 1) \geq p - 1$. The $Q = 1$ branch should bypass the division by $Q - 1$ entirely. Treating positive and negative multiples as unrelated would lose the improved constant in Theorem II.79.3.

Finally, graph energy must be identified with `phiAdditiveCount`. A new order- k theorem should be proved first, and the source's order-eight theorem recovered as a weakened corollary. This preserves existing downstream statements while making stronger constants available to later work.

V.97.5 What would constitute formal verification

The results in this article have not been added to the repository or checked by Lean. Source inspection establishes the existing interfaces, not successful compilation of a modified proof. A completed formalization would include the parameterized finite identities, all injection and threshold lemmas, the final order- k compiler, and a kernel-checked proof of the displayed order-eight bound, with no new axioms or admitted proof holes. The Python regression suite is not a substitute for any of these obligations.

V.98 Further research questions and concrete next steps

Source 24, delivered Section 12.

The questions below are proposed research targets. They are not presented as a list of established famous open problems, nor as claims that the literature has been exhausted. Each is tied to a specific loss or limitation exposed by the proofs.

Question V.98.1 (The actual dense-intersection extremal problem). *For fixed δ , determine the largest fraction $f(\delta)$ such that hypothesis (II.74.1) always supplies a core with 90% of its ordered pairs above threshold $\delta^2 m/2$ and size at least $f(\delta)n$. Is the power δ^3 an artefact of the sample-core construction? Seek matching incidence constructions or a stronger lower bound before making an optimality claim.*

The theorem here gives $f(\delta) \geq \sqrt{C_3}\delta^3$. The scalar sharpness proof does not supply an upper bound for f : its two-point laws may fail to arise from set intersections. A productive first step is to formulate the finite incidence optimization with the positive-semidefinite Gram constraints kept.

Question V.98.2 (Can codegree geometry rescue two samples?). *Does positive semidefiniteness, nonnegative incidence factorization, or control of diagonal entries force a positive two-sample score in a range where the scalar relaxation is negative? Alternatively, can actual incidence families approximate the scalar near-extremizers from Theorem II.75.1?*

A resolution would identify whether the obstruction in (II.75.5) is genuinely combinatorial or only an information-loss barrier. This is a sharper target than simply asking for “better dependent random choice.”

Question V.98.3 (Optimal profile-sensitive sampling). *Given an explicit codegree matrix, how much can one improve the certificate by choosing a nonuniform law ν , a different number of samples, or tuple-dependent thresholds? Can a useful optimization algorithm return a small exact certificate that a proof assistant can verify independently of the search procedure?*

The weighted identity in Theorem II.73.1 is already a correct verification interface. The remaining task is an effective and mathematically justified search strategy, not a new probabilistic existence principle.

Question V.98.4 (Higher-arity scalar envelopes). *For respected arrangements involving r indices, optimize the moment-defect bound in Corollary II.73.2, including the choice of sample number and defect penalty. Which choices minimize the eventual exponent after several structural reductions?*

A natural first case is $r = 3$ with two prescribed badness thresholds. The scalar objective becomes piecewise polynomial, suggesting a finite supporting-line or supporting-polynomial certificate. Its compatibility with actual higher-order intersection tensors remains a separate issue.

Question V.98.5 (The smallest possible vertical-fibre exponent). *Given $|\Gamma - \Gamma| \leq C|\Gamma|$, what is the best universal bound for $|\{y : (0, y) \in k\Gamma - k\Gamma\}|$? Can $2k + 1$ in (II.79.3) be reduced, perhaps under an additional hypothesis on the domain projection or the codomain?*

The first meaningful cases are $k \geq 2$; at $k = 1$ a graph has no nonzero vertical defect at all. Both upper bounds and examples must respect the graph condition. General small-doubling examples without injective first projection do not automatically answer this question.

Question V.98.6 (Sharper extraction from a given symmetric defect set). *For $F \subseteq \mathbb{F}_p$ symmetric and containing zero, what is the largest guaranteed size of a set $P \subseteq \mathbb{F}_p$ satisfying $(kP - kP) \cap F = \{0\}$? Can a nonprogression choice of P improve the extraction fraction beyond the interval-avoidance bound for structured defect sets?*

The proof of Theorem II.79.3 only needs such a set P : a translate of P captures an average fraction $|P|/p$ of the outputs. Its arithmetic-progression construction is explicit but need not be extremal for a particular F .

Question V.98.7 (Mixed-torsion codomains). *Extend the extraction theorem to general finite abelian codomains with quantitative dependence on their invariant factors. Can the prime-cyclic progression method and finite-field hashing be combined without multiplying all the losses from the primary components?*

The obstruction is concrete: the equations $jd = w$ need not have a unique solution when j is not invertible, and a random homomorphism need not send every nonzero defect uniformly to a common target. Any extension should state exactly which torsion parameters enter.

[write] **7 October 2026; sources 54–59.** *Answered by source 59* (Section II.261): exact Freiman restrictions of order k over arbitrary abelian targets, $|B|/(5k)^{\lfloor \log_2(q+1) \rfloor}$ from q sign orbits of vertical defects and $C^{-(2k+1)\log_2(5k)}|B|$ from $|\Gamma - \Gamma| \leq C|\Gamma|$, independently of the target's size, rank and invariant factors, with variants by the defect orders (its Theorem II.264.1 and Corollary II.265.2; item II.8).

Question V.98.8 (End-to-end threshold optimization). *Optimize all structural parameters jointly along the Section 7 pipeline, using the actual energy spectrum rather than a single energy lower bound. How much is gained by combining the full profile certificate (II.76.3), the exact defect count Q_k , and a density-dependent popularity threshold?*

Equation (II.80.5) solves the one-parameter version with coarse energy information. The next version should produce an explicit certificate and compare actual retained sizes, not just formal powers of γ .

Question V.98.9 (Propagation into later sections). *Which later approximate-homomorphism and respected-arrangement lemmas remain valid with their hypotheses unchanged after replacing the Section 7 interfaces by the stronger ones here? After tracing every use, what fully explicit improvement results in the complete 2001 quantitative argument?*

This is a finite but nontrivial dependency-audit problem. Improving one intermediate exponent does not justify replacing every occurrence of a source constant mechanically. The intended outcome should be a complete theorem with all thresholds rechecked, not a heuristic extrapolation to a global record.

Question V.98.10 (A parameterized, kernel-checked structural library). *Can the DRC score identity, graph-walk injection, fibre compression, and quantitative compiler be implemented once over general finite sets and abelian groups, then instantiated both at the original constants and at the improved ones?*

A successful implementation would make future improvements local: new scalar inequalities or new energy inputs could be plugged into unchanged combinatorial proofs. It would also expose exactly which constants are consequences of counting identities and which are chosen for convenience.

V.99 Conclusion

Source 24, delivered Section 13.

The main methodological lesson is to postpone coarse substitution. Keeping the exact codegree distribution allows a cubic certificate that preserves the original dense-intersection threshold. Keeping the actual popular-edge mass cancels four otherwise costly density powers in a walk count. Keeping only the vertical obstruction to Freiman coherence reduces the sumset growth exponent from $4k + 1$, while its symmetry improves the extraction prefactor.

These refinements yield complete local theorems, including the explicit order-eight bound $2^{-295}\gamma^{103}|B_0|$ and the modern-input family (II.80.2). Their scope is precise: proved Section 7 improvements and generalizations, a sharp scalar optimization, and a concrete formalization programme. No unverified propagation through the full Szemerédi argument is needed for any of the conclusions stated here.

V.100 A compact proof and constant ledger

Source 24, delivered Appendix A.

The self-contained order-eight route uses the following exact chain:

$$\begin{aligned}
E(\Gamma_0) &\geq \gamma m^3, \\
\tau &= \gamma/2 \quad (\text{or use the actual energy first}), \\
p &\geq \gamma/(2 - \gamma) \geq \gamma/2, \\
|K| &\geq pm/\sqrt{2}, \\
|\Gamma_1| &\geq |K|/2 \geq 2^{-3}\gamma m, \\
C_1 := \frac{|\Gamma_1 - \Gamma_1|}{|\Gamma_1|} &\leq 2^{11}p^{-2}\tau^{-4} \leq 2^{17}\gamma^{-6}, \\
Q_8 &\leq C_1^{17}, \\
|B| &\geq \frac{|\Gamma_1|}{8Q_8} \geq \frac{2^{-3}\gamma m}{2^3(2^{17}\gamma^{-6})^{17}} \\
&= 2^{-295}\gamma^{103}m.
\end{aligned}$$

When the energy hypothesis is an inequality, either use $\kappa = E(\Gamma_0)/m^3$ throughout and weaken at the end, or retain γ as a lower bound in the popular-energy split. Both give the displayed chain.

For the cubic selector, the arithmetic certificate consists of

$$\begin{aligned} L(s) &= 3(7u/6)^2s - 2(7u/6)^3, \\ s^3 - L(s) &= (s - 7u/6)^2(s + 7u/3), \\ -9(u/2)^3 - L(u/2) &= u^3/108, \\ L(u) &= 49u^3/54. \end{aligned}$$

All constants in these two ledgers are rational except the harmless square root in the intermediate core-size bound. The supplied test program checks the final exponent identities separately from the finite combinatorial examples.

V.101 Package contents and rebuilding

Source 24, delivered Appendix B.

The package contains the article source and PDF, the exact `verify.py` regression program, its recorded `verification_results.json`, a README with scope and build instructions, and a source-provenance note. The article’s bibliography is embedded in this single `.tex` file; no external bibliography processor is required.

From the package directory, the PDF can be rebuilt with

```
pdflatex -interaction=nonstopmode -halt-on-error
gowers_section7_refinements.tex
```

Run the command three times to stabilize the table of contents and cross-references. The line break above is typographic: the command and filename belong on one shell line. The exact finite checks are reproduced with

```
python3 verify.py -output verification_results.json
```

No network access or third-party Python package is required by the test program. A conventional LaTeX installation with the packages listed in the preamble is sufficient for the article.

V.102 Repository integration and verification record

Source 25, delivered Section 9.

V.102.1 Interface crosswalk

A suitable standalone destination is

```
Combinatorics/Ramsey/Research/GowersSzemerédi/arrangement-degeneracy-codimension/.
```

This package does not modify the repository, its ledger, or existing formal proofs. It supplies a separately labelled research manuscript.

Source interface	Result here	Integration point
<code>signedWalshMoment</code> , <code>wals</code> <code>hCharacter</code> <code>arrangementMoment</code>	Lemma II.87.1 Theorem II.88.1	Prove the expansion first, then exhibit its diagonal pivot minor. Split the minimal Walsh degree and the span of the vectors b_i .
<code>GeneralArrangement.IsD</code> <code>egenerate</code>	Theorems II.89.1 and II.90.1	Separate block-parity witnesses; preserve the modular-multiple condition exactly.
<code>degenerateGeneralArrangementCount</code> , <code>lemma_15_4</code>	Theorem F.30.1	Add a refined theorem with $d \geq 2$ and $N > 2d$ prime; retain the old statement for its wider parameter range.
<code>lemma_15_5</code> , <code>arrangementSelectionExponent</code>	Theorem II.92.1	Supply an explicit threshold and exact repetition count before deciding whether to weaken them to the existing statement.

For a first formalization, using $\mathbb{Z}/N\mathbb{Z}$ with N prime would minimize changes to the existing definitions. The extension to an abstract finite field can then be a generalization of the new result. The geometric codimension statement need not be formalized to obtain the finite counting bound: pivot-coordinate bijections and finite linear-map fibre cardinalities suffice.

The hyperplane enumeration should be a separate finite combinatorics lemma. In particular, the quotient by $a \sim \lambda a + \mu w$ must not be replaced by the simpler but incorrect quotient $a \sim -a$. The two-value classes have six representatives, not two. Likewise, the constant cross-section fibre in Theorem II.88.1 must not be omitted; it is the source of the exact $(q-1)q^{-2d}$ term.

No Lean source is included with placeholder proofs. This avoids giving unverified code the appearance of a checked formal result. A natural implementation order is Walsh expansion, pivot ranks, the two-row fibre count, projective classes, union bounds, and finally selection.

V.102.2 Exact computations supplied with the manuscript

The script `code/verify.py` uses only Python’s standard library. Its saved output is `data/verification.json`. The executed checks are as follows.

- (1) Direct integer-projective and finite-field enumeration of the hyperplane classes for $d = 1, \dots, 5$, matching 0, 4, 51, 484, 4225; evaluation of the closed formula through $d = 9$.
- (2) An exact characteristic-5 check for $d = 3$, obtaining 57 rather than 51 hyperplanes.
- (3) 1550 instances of the Walsh moment identity, and 315 seeded minimal-layer rank tests, with cube dimensions up to six.
- (4) All 3240 non-block-parity ternary sign classes for $k = 1, d = 2$ over each of $\mathbb{F}_5$ and $\mathbb{F}_7$. For each class the full base-coordinate count is obtained from exact ranks of the two affine equations, over all cross-sections. The maximal probabilities are respectively 29/625 and 55/2401, exactly the bound of Theorem II.88.1.

- (5) Exact counts for the sharp witness with $(d, q) = (1, 5), (2, 5), (2, 7), (3, 7)$, and rational checks of the one-factor Riesz coefficients for $n = 8, 16, 32, 64$.

All these checks passed in the recorded run. They test finite cases, normalizations and algebraic identities; they do not replace the general proofs. In particular, the script does not enumerate every full arrangement, does not simulate the selection theorem, and does not run Lean. The source audit and README in the archive repeat these boundaries for repository intake.

[write] **6 October 2026; sources 21–27.** *Checked at HEAD 2dfe46f44.* Every name in the table exists in `Sections14_15.lean`: `walshCharacter` (line 310), `signedWalshMoment` (line 329), `arrangementMoment` (line 346), `GeneralArrangement.IsDegenerate` (line 354), `degenerateGeneralArrangementCount` (line 364), `lemma_15_4` (line 402), `arrangementSelectionExponent` (line 409) and `lemma_15_5` (line 414); the file has the same blob as at this source’s pin. The predicate requires all moments, including the empty monomial and the last coordinate, and excludes modular multiples, as the definition here does. The suggested separate directory was declined: the manuscript is printed in this report.

V.103 Further research questions

Source 25, delivered Section 10.

Question V.103.1 (The exact second-order term). *For fixed k, d , determine whether a field-independent coefficient $C_{k,d}$ exists such that*

$$Z_{k,d}(q) = (k + H_d)q^{D-1} + C_{k,d}q^{D-2} + o(q^{D-2})$$

over sufficiently large characteristics, and compute it in the first nontrivial cases. Theorem II.90.1 proves that non-parity witnesses already contribute at this order. Intersections of the known hyperplanes and overlaps among codimension-two witness loci must both be counted.

Question V.103.2 (Compressing the second-order constant). *Replace $U_{k,d} = (3^{2d2^k} - 3^{2d})/2$ by a substantially smaller explicit constant. Many ternary patterns have identical or nested moment loci. The minimal-layer rank $k - s$ suggests separating witnesses by Walsh support before applying the union bound, rather than assigning every remaining witness the same q^{-2} cost.*

Question V.103.3 (The full cross-section arrangement). *Determine the intersection lattice and characteristic polynomial of $\mathcal{H}_d$. For $d = 2$, the exact union count is (II.90.5); for larger d , a structural description would replace the Bonferroni estimate with an exact formula for $v_d(q)$. The classical finite-field method for hyperplane arrangements [98] is a natural framework, but a full analysis of this specific balanced ternary quotient is not supplied here.*

Question V.103.4 (Small-characteristic classification). *Find a closed formula for $|\mathcal{H}_d(\mathbb{F}_q)|$ when $3 < p \leq 2d$, and classify the additional components in characteristics 2 and 3. Example II.89.2 shows that replacing the characteristic assumption by “ q sufficiently large” is not valid. A satisfactory answer should identify the modular balancing terms and the altered affine orbits, rather than merely give numerical tables.*

Question V.103.5 (Efficient detection of degeneracy). *Given an arrangement, detect a nontrivial ternary moment relation without scanning 3^n patterns. The present classification supplies a cheap first-stage test using $k + H_d$ hyperplanes; a useful second stage would exploit the structured Walsh moment matrix and return a verifiable witness. The difficulty is ternary dependence, not ordinary linear kernel computation, since the moment matrix has many more columns than rows in the applications.*

Question V.103.6 (A selection mechanism with a larger advantage). *Can a different nonnegative selection kernel distinguish respected and non-respected arrangements by more than the ratio $1 + 2^{1-n}$, while retaining manageable degeneracy conditions and pointwise probabilities in $[0, 1]$? Better exceptional-set counting alone does not remove the exponential loss in (II.92.12). Higher Fourier support may improve the advantage but introduces coefficient relations beyond ternary ones.*

Question V.103.7 (Uniform extension to other defining relations). *Replace $w = (1^d, -1^d)$ by an arbitrary nonzero coefficient vector w with $\sum_j w_j = 0$. Which support and coefficient conditions force every non-parity witness to have codimension at least two? The rank argument indicates the role of the support of w , but the exact projective hyperplane count and the exceptional low-support cases require a separate treatment.*

Question V.103.8 (End-to-end quantitative propagation). *Integrate the explicit selection threshold into the corrected multilinear-piece and final density-increment arguments in ProveIt. Track both the field-size threshold and the retained arrangement density. Only such an end-to-end calculation could justify a claim about an improved final Szemerédi or van der Waerden bound. The current manuscript deliberately stops short of that claim.*

V.104 Conclusion

Source 25, delivered Section 11.

The full arrangement moment system is more rigid than the one-polynomial estimate used to count its exceptions. A lowest nonzero Walsh layer gives an explicit diagonal linear minor. At the sole remaining adjacent layer, the last-coordinate moments force a second independent equation, except on already visible cross-section degeneracies. As a result, all codimension-one exceptions are classified and their exact number is $k + H_d$, with H_d given by a balanced ternary projective count.

This yields an asymptotically optimal local replacement for Lemma 15.4, a finite-field and arbitrary- d purification theorem with explicit parameters, and a concrete reduction of the ambient-size cost of the exceptional set. The unresolved issues are correspondingly more focused: second-order geometry, compression of the residual witness family, and improvement or replacement of the Riesz-product selection advantage.

V.105 Source comparison and remaining questions

Source 26, delivered Section 5.

V.105.1 The precise repository interface

The mathematical source pin is commit

`c94998a5db18ba63c815a3bf6c3877a0fde91930`.

At that pin, `Section05.lean`, lines 242–259, defines the proposition `corollary_5_8` [116]. It assumes disjoint two-comparable proper modular progressions covering A , local degree-at-most- k polynomial correlations, and $\alpha > 0$. Its length conclusion is

$$\frac{(N/M)^{1/K}}{8}, \quad K = (k!)^2 2^{(k+1)^2},$$

and its density conclusion is $\delta + \alpha/16$. The original paper prints an $\alpha/8$ gain in Corollary 5.8, on printed page 493 (PDF page 29). The source comment records that this conclusion was corrected to the $\alpha/16$ actually supported by the proof. We compare with this pinned corrected interface rather than silently switching between the two constants.

The definition of `PolynomialOn 1 Finset.univ` is an existential coefficient sum over `Fin 2`: it expands to $c_0 + c_1x$ [117]. Thus the affine hypothesis here is exactly the degree-one source condition, not an extra linearization assumption. At $k = 1$, $K = 16$. For every modulus $N \geq 2$, Corollary I.41.3 replaces the source’s length $(N/M)^{1/16}/8$ by $(N/M)^{1/3}/8$ and the gain $\alpha/16$ by $\alpha/3$, with no added scale premise. Nonempty parents are automatic in that source specialization: positive correlation implies $A \neq \emptyset$, and comparability would force all parents to be empty if any one were empty. Their disjointness then gives $M \leq N$, so $N/M \geq 1$ and the larger exponent indeed strengthens the old length bound throughout the nonvacuous range.

The separate source file `Proofs05Corollary58.lean` defines a scale-qualified repair with a common target and lower bounds ensuring that the needed recurrence and phase-freezing estimates apply [118]. Its theorem derives that repaired proposition from `corollary_5_6`. Those added scale conditions are not premises of the displayed catalogue predicate. The argument here supplies a written, threshold-free proof of the full degree-one specialization; it is not a Lean implementation of that specialization.

[write] **6 October 2026; sources 21–27.** *Checked at HEAD 2dfe46f44.* The pinned lines 242–259 of `Section05.lean` are unchanged (same blob), and `corollary_5_8` is still open in the ledger; `PolynomialOn` is at `Definitions.lean`, line 222, and `corollary_5_8_with_scale_holds` at `Proofs05FullPartition.lean`, line 94. The “full degree-one specialization” holds trivially when $N/M \leq 2^{48}$ (note in Section F.31).

The remaining source modulus $N = 1$ is vacuous: its balanced function is identically zero, so the positive-correlation hypothesis cannot hold. Thus the written result covers every modulus in the source’s degree-one specialization. Higher-degree phases are outside the present theorems. The output is proper modular geometry, as encoded in the pinned Lean definitions; it need not satisfy every ordinary-integer geometric convention of the original paper.

V.105.2 The classical mechanisms and the limited novelty claim

Gowers’s proof contains the Dirichlet input and the square-root modular-to-integer decomposition: see Lemma 5.4, the one-axis Lemma 5.12, and its whole-partition consequence Lemma 5.13 [1]. The product almost-cover in Appendix I.44 is an elementary explicit wrapper around this type of construction.

Zhao’s Lemma 6.4.5 partitions an interval into phase-flat progressions on the cube-root scale under an explicit size threshold; Lemma 6.4.6 uses such a partition in a Roth density increment [100]. This provides a direct reason not to portray the cube-root exponent or phase partitioning as new. Lewko’s 2026 preprint uses moments of progression sums in Section 5 and extraction of a genuine integer progression in Section 6 [101]. It likewise provides precedent for these methods; it is not invoked as a proof of our exact local-cover statements.

The contribution established in this report is the explicit threshold-free local-cover combination, its exact-mass formulation and correlation-dependent tradeoff, relative to the pinned repository interface. A comprehensive literature-priority investigation has not been completed. The numerical constants are convenient and are not claimed optimal.

V.105.3 Questions that remain open here

1. Can the explicit $(N/8)^{1/3}$ ceiling in the strongest exact-mass formulation be relaxed? It arises from our sufficient condition $N \geq L^3$ for removing nonproper directions; it does not restrict the simpler normalized-cover or comparable-parent conclusions.
2. Can correlations be used more efficiently to obtain an α -independent square-root guarantee? Theorem I.41.2 proves only an α -dependent square-root alternative.
3. Which exact constants and gains are optimal for this local-cover interface? The ceiling/floor losses, $\pi < 4$, and positive-mass averaging leave deliberate slack.
4. What is a threshold-free higher-degree replacement for the phase-flat refinement, and how would its small-correlation regime be handled? No such result is supplied here.
5. Can the elementary pair-counting, residue partition, and signed-mass lemmas be formalized and then composed into the precise all-modulus $k = 1$ predicate? No Lean theorem was produced or kernel-checked in this report.
6. Can the integer-box almost-cover be connected to a source-faithful exact-partition and base-loss interface? Appendix I.44.1 identifies the two distinct obligations.

V.106 The exact companion and reproducibility

Source 26, delivered Section 6.

The Python companion uses only the standard library, integer arithmetic, and exact rational arithmetic for its mathematical checks. Its role is diagnostic: finite tests can expose a wrong midpoint, a pair-counting failure, an omitted endpoint, or an invalid rounding convention. They cannot prove the arbitrary-parameter theorems, which are proved in Sections I.42–I.43.

The companion checks cover ordered modular pair counts and the exact variance identity in its permitted range, all-direction coset covariance and bad-direction removal, elementary Dirichlet residue cells and lossless balanced chunking, the exact cover-mass identities, and a regression distinguishing the reduced and unreduced midpoints. The command-line bound is finite and capped. The optional larger run changes the finite coverage, not the proof status. The source archive's README gives the exact commands and coverage labels printed by the final companion.

The offline builder compiles the supplied trusted TeX in fresh external staging, runs tests and diagnostics in normal and optimized Python modes, rejects inventory mismatches, and writes a deterministic ZIP with an exact public-file manifest. It does not download dependencies or require network access. Reproducing the PDF byte-for-byte requires the same TeX toolchain and fonts. The file-handling checks and disabled shell escape are useful safeguards for this trusted-source workflow; they are not a sandbox guarantee for hostile TeX or a concurrently malicious filesystem. The archive includes all source and tests required for its stated checks.

V.107 Verification and repository integration

Source 27, delivered Section 12.

V.107.1 What the supplied computation verifies

The package contains a single main verifier, `code/verify.py`, and its exact rational data, `code/exact_certificate.json`. The default verifier uses only the Python standard library. It reconstructs the certificate from scratch and checks that it agrees with the stored file. It verifies the 484/459/15 class inventory, every line and point incidence, the 7296 rank certificates, and the exceptional-prime set $\{2, 3, 5, 7\}$. It also enumerates the balanced profile classes for lengths 4, 6, 8 in characteristics 2, 3, 5, 7, 11, checking all fifteen instances of Theorem II.95.1.

A deterministic additional check constructs all 170 large-characteristic core profiles in the cases $(k, d) = (1, 2), (2, 2), (3, 2), (4, 2)$, chooses bases satisfying (II.97.5), and verifies every squarefree moment exactly modulo the chosen prime. These checks test the signs and normalizations in the Boolean-core formula; they are not substitutes for its proof.

The optional `--subset` mode uses NumPy to perform the independent exhaustive tests in Table II.100.2. It does not import or use the rational line certificate. The command

```
python3 code/verify.py -subset -write
```

regenerates the stored verification report. Without NumPy, omit `--subset`. The optional `--minors` flag, together with `--write`, writes the individual nonzero-minor witnesses; these are reconstructible and need not be stored to reproduce the proof.

No floating-point computation is used to prove a count, rank, or incidence. No numerical experiments are presented as evidence of novelty. The code has been executed successfully for the distributed report; its results are recorded in `verification.json`.

V.107.2 Source correspondence

The following map uses names actually present in the inspected statement file [120]. It proposes new additions; it does not claim that any such additions already exist.

Table V.107.1: Correspondence with the existing Section 15 interface.

Existing name or mathematical object	Proposed use
<code>GeneralArrangement</code>	Preserve the labelled side, bases, and cross-section parameters. Specialize the present field to a prime $\mathbb{Z}\text{Mod}N$ first.
<code>arrangementParityCoefficient</code> <code>arrangementMoment</code>	Identify it with $e_{\varepsilon,j}$ and prove all its moments vanish. Prove the Boolean expansion and the paired moment identities. Do not delete any subset index.
<code>GeneralArrangement.IsDegenerate</code>	Retain the ternary condition and modular-multiple exclusion; build the quotient-direction counting lemmas around this existing definition.
<code>degenerateGeneralArrangementCount</code>	Multiply the probability bounds by N^D to obtain integer-count statements, then cast to the existing real-valued bound style.

Existing name or mathematical object	Proposed use
<code>lemma_15_4</code>	Add a sharper companion theorem, preserving the original statement and proof ledger rather than silently changing them.
<code>lemma_15_5</code>	Use the new bound only at the place where an absolute exceptional-arrangement budget is needed; rerun that threshold calculation separately.

V.107.3 A realistic formalization order

A useful first milestone is Theorem II.96.3 over odd prime fields, because it strengthens the existing interface without requiring the most detailed core classification. Its dependencies are: intrinsic moment vanishing; Boolean inversion; codimension-one counts in U ; the elementary rank lemma; and a finite union bound. The ternary affine-orbit formulas can initially be replaced by finite cardinalities of explicitly defined sets, then proved as a separate arithmetic-combinatorial layer.

A second milestone is Theorem II.97.8: formalize the minimal Boolean degree split, the affine core form, the width-two classification, and the rank-four intersection statement. The all-characteristic version should use a finite-field abstraction rather than force the extension-field argument through $\mathbf{ZModN}$ for composite N , which is generally not a field.

The exact theorem has a distinct final dependency: a verified finite certificate checker for rational row spaces and reduction modulo permitted primes. It should not be imported as an unproved numeric fact. A modest initial target is to formalize (F.32.10), whose proof is entirely symbolic and independent of that certificate.

V.107.4 Provenance and status

The source audit identifies the inspected repository blobs and the printed paper pages. These are snapshots, not a claim that a rapidly changing `main` branch will retain those contents. The research directory already contains numerous local refinements; the present article focuses on the degeneracy count rather than reproducing phase partitions, additive cube expansions, or random-selection redesigns. The targeted review did not establish a complete literature priority result. The mathematical statements here should be treated as proposed research contributions with proofs supplied for review.

No GitHub write, source replacement, or claim of formal completion is made by this delivery. The suggested repository location is

`Combinatorics/Ramsey/Research/GowersSzemerédi/
arrangement-degeneracy/`

with later integration into a merged report at the maintainer’s discretion.

V.108 Further research questions

Source 27, delivered Section 13.

The following questions are deliberately more specific than a request to improve the global theorem. Some are finite classification problems suitable for immediate computation; others require genuinely new structural input. They are proposed questions of this article, not claims that each has been previously advertised as an open problem in the literature.

Research question V.108.1 (The cross-section intersection statistic). Find a closed form, recurrence, or efficient exact algorithm for $b_{2,d,p}$ in (II.99.1), particularly in the stable large-characteristic regime. This would turn Corollary II.99.2 into a numerical second-order formula for every d , including the $d = 8$ case of the selection step.

A first computational target is $d = 3$, where there are 51 hyperplanes in large characteristic. Enumerating their pair intersections is much smaller than enumerating all arrangement coefficient arrays. The main issue is recognizing coincident rank-two spans, not sampling arrangements.

Research question V.108.2 (A third-order structural theorem). Classify the non-core coefficient directions whose zero sets on the regular region contribute at order q^{-3} . Can the residual term $(R_{k,d} - S_{k,d})q^{-3}$ be replaced by an explicit third-order family plus an $O(q^{-4})$ remainder?

The proof already distinguishes directions with a new cross-section constraint from those forcing two separate rank-two base systems. The latter are automatically fourth-order. Exploiting this distinction before using the final union bound should reduce the next-stage inventory.

Research question V.108.3 (Higher Boolean degree and increasing codimension). For a direction of minimal Boolean coefficient degree r , how many independent base-coordinate systems can be extracted simultaneously? Is there a sharp codimension bound in terms of $k - r$ and the ranks of the vectors c_A modulo s ?

The rank-four argument uses just two coordinates outside one minimal set. A broader theorem would need to control overlaps among right sides and cross-section constraints, rather than repeatedly invoke independence without checking it.

Research question V.108.4 (Intermediate characteristics). Determine the exact number of core directions for $5 \leq p \leq 2k + 2$ and the corresponding intersection pattern. The cases $p = 2, 3$ are answered by Theorem II.99.3, and the width argument already supplies the large-characteristic core count for $p > 2k + 2$. What is the sharp threshold, and which smaller characteristics have genuinely different core families?

The remaining gap is to classify profiles below the proved no-wrap threshold, not merely to replace $p > m$ by $p > 2k + 2$. For larger alphabets, however, new affine symmetries and wider Boolean profiles must be accounted for.

Research question V.108.5 (Exact counts in fixed small characteristic). Find exact formulas for $\Delta_{1,2}(p^a)$ in characteristics 2, 3, 5, 7. The first two coefficients in characteristics 2, 3 are already determined here; the large-characteristic polynomial fails at $q = 5, 7$.

The same line-enumeration method can be performed directly over $\mathbb{F}_p$ and then lifted to extensions. In characteristic two, the subset-sum test must be reformulated because the two integer signs have the same residue.

Research question V.108.6 (Beyond eight vertices). For $(k, d) = (1, 3)$ or $(2, 2)$, is the degeneracy count eventually a polynomial in q outside finitely many characteristics? If not, what arithmetic variation occurs in the lower-dimensional residual loci?

Only the hyperplane part is automatically a polynomial count in fixed characteristic. The full higher-degree moment varieties need not inherit that property without further analysis, so polynomiality is a question, not an inference from the eight-vertex example.

Research question V.108.7 (Minimal moment families). Which subfamilies of squarefree moments retain the same first-order or second-order degeneracy asymptotics? Can one preserve the cubic-remainder bound while discarding a substantial fraction of the moment conditions?

The paired moments used in Lemma II.94.2 identify a necessary organizing principle. A universal fixed subfamily would have to work for every possible minimal Boolean coefficient set; the weaker definition mentioned after Gowers’s Lemma 15.4 is not automatically sufficient.

Research question V.108.8 (Efficient degeneracy testing). Can the quotient-profile structure yield a deterministic test significantly faster than enumerating 3^m relations or 2^m subset sums, while producing a short certificate of either degeneracy or nondegeneracy?

The exact low-dimensional example suggests preprocessing by cross-section hyperplanes and then by the core family. A useful algorithm must also handle the remaining coefficient directions without an exponential worst-case scan disguised as preprocessing.

Research question V.108.9 (Relative distribution rather than ambient bounds). For arrangement collections arising from a dense set or a product-property hypothesis, can the hyperplane contribution be bounded by their actual cross-section distribution instead of the ambient L_d/q union bound?

Proposition II.101.1 loses a factor equal to the reciprocal of collection density. Avoiding it would require structural information about the collection, not just its cardinality. A weighted version could be directly relevant to selection arguments.

Research question V.108.10 (The single-pair regime). Classify the dominant degeneracy loci for $d = 1$ and arbitrary k . Does the common final coordinate reduce the problem to a useful arrangement of relations among two translated cubes?

Proposition II.102.1 settles only $k = 1$. The rank collapse is systematic and should be treated as a separate geometry rather than forced into the $d \geq 2$ proof.

Research question V.108.11 (A complete Section 15 threshold audit). Substitute the cubic bound into the actual corrected Lemma 15.5 proof and propagate the resulting threshold through Lemma 15.6 and the later inverse steps. Which local improvement survives, and which other parameter becomes dominant?

This is the precise route from the present local result toward a global quantitative consequence. It needs a complete ledger of dependencies and is not settled by the numerical comparisons in Table II.101.1.

Research question V.108.12 (Formal proof and verified finite certificates). Formalize the sharp first-order theorem, then the second-order theorem for $d = 2$, and finally the exact eight-vertex count with a small verified rational-incidence checker. Which combinatorial and finite-dimensional linear-algebra interfaces can be reused in other parts of Gowers’s proof?

The separation between symbolic general proofs and a finite exact certificate makes this an incremental task: useful theorems can be checked before the final enumeration infrastructure is complete.

V.109 Conclusion

Source 27, delivered Section 14.

The large coefficient in Gowers’s degeneracy bound is not the true leading coefficient of the exceptional-set count. After coefficient arrays are identified modulo the intrinsic relation, the bad set splits into three quantitatively different parts: explicit hyperplanes, a classified family of rank-two affine systems, and a residual set costing at least three powers of the field size. The resulting estimates are sharp to first order in all characteristics and to second order once the cross-section hyperplane statistic is known. For two pairs of cross sections that statistic is explicit, and in the smallest nontrivial case the entire count is an exact polynomial.

These results answer concrete local questions about the proof: what actually creates the leading degeneracy, how much of the ternary union bound is avoidable, and which additional moment constraints supply the extra savings. They also leave a clear boundary between established local proofs, a finite certified specialization, and the further work needed for new global consequences.

V.110 Reproduction checklist and file inventory

Source 27, delivered Appendix A.

The archive contains the following files, all under one top-level directory:

File	Purpose
<code>article.tex</code>	Complete self-contained LaTeX source
<code>article.pdf</code>	Compiled article
<code>README.md</code>	Results, scope, and reproduction instructions
<code>SOURCE_AUDIT.md</code>	Primary sources and repository snapshot identifiers
<code>FORMALIZATION.md</code>	Proposed Lean dependency plan
<code>Makefile</code>	Build and verification targets
<code>code/verify.py</code>	Exact reconstruction and independent checks
<code>code/exact_certificate.json</code>	Rational lines, points, and incidences
<code>verification.json</code>	Results of the executed checks

Run `makepdf` with a standard TeX installation containing the packages listed in the preamble. Run `makeverify` for the standard-library checks, or `makeverify-full` for the additional NumPy subset-sum checks. The article uses no external figures, fonts, bibliography database, shell escape, or network requests. Its source can therefore be compiled independently of the verifier.

V.111 An explicit interface for further development

Source 28, delivered Section 9.

Delivered part heading of source 28: **Verification, integration, and further research**

V.111.1 What a downstream theorem can use

The first part produces four distinct kinds of information. Keeping their roles separate makes an integration into the existing argument more transparent.

1. **A prescribed spectrum.** The character set is obtained from $g(s) = |X_s|$, with the threshold in equation (F.33.2). No label information is used to choose it. A downstream step can therefore synchronize spectra using domain data.
2. **A large set of valid labels.** The set Y has the explicit relative size $\alpha p(\alpha)/20$. It is selected after seeing ϕ . Its projection to G need not have that relative size, since distinct points in a fibre can have multiplicity.
3. **An exact map on a neighborhood.** The identity equation (F.33.7) holds for all eligible pairs in Y . The local map respects order-seven relations on its entire Bohr domain, not only on the differences realized by pairs of Y .
4. **Quantitative geometric control.** Both the frequency count and the radius are explicit. The resulting volume estimate is equation (II.127.2). For an existing application that fixes the old frequency set, corollary II.126.2 supplies a compatible restricted map with the same retained fraction and radius.

These properties should be passed as separate fields in a theorem interface. A mere assertion that a large subset admits some additive model would discard the original-spectrum information. Conversely, requiring a global homomorphism $G \rightarrow H$ would strengthen the conclusion beyond what the proof establishes.

V.111.2 A proposed formalization sequence

The following names are suggestions for new modules, not existing or compiled files in this deliverable. The current article is a mathematical and reproducibility supplement; it contains no Lean code.

Table V.111.1: Suggested proof interfaces and their dependency order.

Suggested module	Mathematical content to formalize
<code>WeightedAnchor</code>	Finite energy bounds, local error transfer, and the single weighted selection of an anchor in section II.123.
<code>LogWindowSelection</code>	The two center intervals, their clipped inclusion probabilities, the weighted boundary estimate, and equation (II.124.8).
<code>AnchoredModalRepair</code>	Comparison between the $b_d(w)$ -weighted error and uniform direction error, existence of a unique modal increment, and the row restriction.
<code>StoppedPathRigidity</code>	Subprobability transport with killing on exits and wrong labels; shared endpoint coupling; the order-fifteen conclusion.
<code>OriginalSpectrumModel</code>	Fourfold smoothing, radius selection before the anchor is chosen, the final slice, and the order-seven local model.

Suggested module	Mathematical content to formalize
<code>UFourDualCosine</code>	The integer sign counts 222, 111, 27, 1, including the separate order-three quotient calculation.
<code>HigherUniformityGap</code>	The concentration inequalities, rational threshold comparison, two-harmonic certificate, and positive first-variation recurrence.

The logarithmic-window argument has a finite alternative to an abstract probability-space formalization. Each point contributes finitely many center endpoints: its two outer endpoints and its two inner endpoints in each coordinate. Together with the endpoints of I_q, I_R , they partition the center rectangle into finitely many open rectangles. Membership in W , membership in its boundary, and the selection functional are constant on each such rectangle. Their weighted finite average is the integral in equation (II.124.8). A rectangle with value at least the average yields the chosen window. Endpoints have zero area and can be handled independently.

This reduction also suggests a constructive, though not optimized, selection algorithm. Compute the anchor statistics, sort the center endpoints, evaluate one interior center in each nonempty rectangle, and choose a maximum of the functional. Real logarithm comparisons would still need certified handling for an exact implementation; the present program tests the finite mechanisms and does not implement a certified end-to-end model finder.

V.111.3 The existing formalization remains a separate object

The inspected status ledger explicitly distinguishes established theorems, conditional results, migrated statements, and outstanding obligations [131]. Several prerequisites in Sections 13, 16, and 18 remain open at the pinned snapshot. In particular, the source records the need for additional bilinear extraction, a closing quantitative comparison in the multilinear stage, and sufficiently strong phase removal and iteration thresholds.

Some of those issues concern the statement being formalized rather than an isolated numerical loss. Examples recorded by the repository include integer representatives of box axes, shared partition requirements, and a transcription’s closing comparison whose proposed direction has already been disproved. None can be discharged by inserting a better Section 10 constant. A proposed integration should first state the exact downstream predicate and then verify that the local model satisfies it.

V.112 Exact computations and the scope of verification

Source 28, delivered Section 10.

V.112.1 Finite checks of the selection and transport mechanisms

The script `code/verify_selection.py` uses integers and exact rational arithmetic for its main checks. The fixed seed is 20261006; the output records the examples and all summary counts. Its coverage includes:

- 648 exact window-membership cases, including 200 exact core inclusion probabilities and 192 clipped noncore cases;

- 10 368 boundary-probability comparisons, including 6 272 positive boundaries;
- exact integration of the weighted selection functional over 1 483 constant-membership cells from four abstract examples;
- 14 376 shifted-interior checks in three cyclic integer fibre profiles, using 84 proper windows, with 168 actual exits before the boundary is removed;
- all 314 721 ordered pairs of the 561 probability distributions on three symbols with denominator 32, testing the modal-agreement implication;
- 228 coupled two-path checks, including positive label error and mass leaving the selected set;
- stopped-prefix and longer common-endpoint checks that retain unnormalized surviving mass, together with exact verification of the order-fifteen inequality.

For example, the deliberately conservative long-path budget used in the manuscript is the exact fraction

$$\left(30 + 210 \cdot 2^{-27}\right) \left(\frac{1}{31} + 2^{-21}\right) = \frac{4222187282107575}{4362862139015168} < 1.$$

The script also checks the algebraic parameter identities against a rational enclosure of $\log 2$ of width $1/884569308258492$.

A separately labelled numerical section uses actual natural logarithms in three additional profiles, totaling 1 205 cells. The maximum discrepancy between the two evaluations of the expected functional is about $1.49 \cdot 10^{-11}$, within the stated $2 \cdot 10^{-9}$ tolerance. These examples need not satisfy the tiny final invariance parameter. They check the window calculation, not the full approximate-homomorphism theorem. Abstract examples and physical fibre examples are also identified separately in the output.

V.112.2 Exact norm certificates

The script `code/verify_u4.py` uses only integers and rational numbers. It reproduces the following identities and comparisons:

1. the full-cube cosine count and all three pairs of dual coefficients;
2. the order-three count and its phase-dependent minimum;
3. every coefficient of the two-harmonic polynomial, by combining two three-dimensional faces;
4. the five-point example, by exact derivative energies;
5. the seven-point example, both by modular coefficient counting and by direct cube summation in a cubic number field;
6. the rational split at $159/160$, including both alternatives in the universal gap;
7. the exact value of the witness at $b = -1/7$, and explicit one-third-harmonic assignments in dimensions 4 through 8.

The polynomial identity is reproducible by enumerating $4^8 = 65\,536$ choices on one three-face. Group choices by their moment vector and number of third harmonics, then convolve the count polynomial for each vector with the one for its negative. This computes equation (III.70.20) directly. The generic and small-order computations are intentionally separate: replacing equality of integers by congruence changes the answer.

The rational comparison in the universal bound can be checked without decimal approximations. One of its final steps is precisely

$$111 \cdot 3200 \cdot 5147^{16} > 104329 \cdot 5550^{16}.$$

The large integers in the witness are likewise exact. The plot in figure III.70.1 has an explanatory role and is not used to infer the sign of a derivative or locate a global optimum.

V.112.3 Independent review and its limits

The selection proof, the averaged repair, and the higher-order norm argument were separately checked by other reasoning agents during the preparation of this report. This review caught interface details, including nonemptiness hypotheses and the need to describe stopped-path marginals as subprobability laws. It also independently reproduced dual-function counts for character orders 3, 5, 7, 9, 11.

Such review is useful evidence of internal consistency. It is not an external referee report, a proof-assistant certificate, or an exhaustive search for counterexamples. The general theorems stand on their written proofs. The package records exact computations to make sensitive finite steps easier to assess and repeat.

V.113 Further research questions and topics

Source 28, delivered Section 11.

The most useful next projects are those that separate a genuine structural obstruction from a loss specific to the present proof. The following questions are proposed in that spirit.

V.113.1 Selection and local structure

Research question V.113.1 (Are both logarithmic losses necessary?). Can the retained fraction in theorem F.33.1 be improved to $\Omega(\alpha/\log(1/\alpha))$, or even $\Omega(\alpha)$, while preserving an original-domain spectrum with comparable size and radius?

There are two independently selected scales here, for q and R . An obstruction for arbitrary weighted profiles would not automatically be an obstruction for the profiles arising from an additive convolution. A useful first step is to construct or rule out domains whose anchored energy is spread nearly uniformly over both logarithmic coordinates. Correlations forced by the convolution may reduce the number of scales that need to be selected.

Research question V.113.2 (Can the defining frequency count be reduced?). Can one replace $O(\alpha^{-3})$ in the local-model theorem by a smaller density power, or by a bound on an independent generating set of frequencies with a controlled loss of radius?

The present count is Parseval's bound α/λ^2 . The main technical requirement is uniform control of the fourfold convolution under the resulting directions. A spectrum compression argument must record how character combinations change the Bohr radius. Counting independent generators and counting all coefficients above a threshold are different tasks, with different downstream interfaces.

Research question V.113.3 (A general defect–mass–radius trade-off). What is the best simultaneous dependence on an arbitrary order-eight defect ε , retained fraction, and local radius?

The proof fixes $\varepsilon = 2^{-21}$ to keep one readable set of constants. Its intermediate inequalities allow parameters such as the modal threshold, window aspect ratios, direction-hole budget, and Bohr-overlap budget to vary. Optimizing those parameters should produce a family of explicit sufficient trade-offs. Whether such a family is sharp would require examples with genuinely inconsistent labels, not just adverse fibre sizes.

Research question V.113.4 (The exact output order). With the same input defect, can the local model be made to preserve order eight, or arbitrary fixed order, without materially weakening its mass and radius?

The current argument obtains order fifteen on D , then order seven on $D - D$. The restriction to fifteen is the displayed path budget; it is not an impossibility theorem. Better use of correlated label errors, a stronger modal threshold, or a direct relation argument on the difference set may improve the conclusion.

Research question V.113.5 (How much input order is needed?). Which conclusions survive if the order-eight assumption is replaced by order four or order two, while the defining spectrum must still come from the original domain?

The exponent eight enters through the fourfold difference domain, whose fibres have a nonnegative fourth-power Fourier transform. Reducing the input order calls for a different smoothing device or a direct argument that retains the same original-spectrum control. An arbitrary structured subset theorem would not resolve this more specific requirement.

Research question V.113.6 (A certified finite construction). Can the logarithmic selection and modal repair be implemented with certified arithmetic in a running time substantially better than enumerating all relations and all center rectangles?

The existence proof gives a finite search once the statistics are available. Useful algorithmic improvements would include implicit weighted rectangle maximization, sparse convolution when the group permits it, and exact decisions at logarithmic endpoints. The output should include a certificate of every required fibre and relation identity, rather than only a candidate model.

V.113.2 Higher-order norm geometry

Research question V.113.7 (The sharp fourth-order constant). Determine c_4^{odd} , or substantially narrow the interval in equation (III.70.4).

The interval is still wide. On the upper-bound side, the present proof compresses all secondary Fourier coefficients into a single remainder mass. Keeping separate information at χ^3 and χ^5 , or using the dual function of a two-harmonic trial function, could give stronger constraints. On the constructive side, additional phase-aligned odd harmonics may improve the explicit example.

Research question V.113.8 (Do bounded spectral profiles suffice?). Is the fourth-order supremum attained by a function with a uniformly bounded number of Fourier coefficients? Is a two-harmonic family ever globally optimal?

The first variation rules out one character pair, but gives no support bound for a maximizer. One can first investigate the exact finite polynomials for three harmonics and their stationarity conditions. A proof that all further harmonics are unfavorable would be a separate and more demanding step.

Research question V.113.9 (Small torsion versus large character order). How does the optimal constant depend on the allowed character orders, and does a fixed finite cyclic group attain the unrestricted supremum?

The order-three branch of the upper bound and the seven-point polynomial show why congruence aliases must be kept visible. One may define suprema restricted to a fixed exponent, to cyclic groups, or to groups with no characters of small order, then compare them. A limiting Fourier profile on the circle could behave differently from every fixed finite model.

Research question V.113.10 (The asymptotic higher-order constants). What is the behavior of c_d^{odd} as $d \rightarrow \infty$?

[write] **7 October 2026; sources 60–69.** *Answered by source 63:* $c_d^{\text{odd}} \rightarrow 1/3$, with $c_d^{\text{odd}} - 1/3 \leq 10 \cdot 2^{-\lfloor d/8 \rfloor}$ for $d \geq 80$; the optimal rate is open (its Research question V.237.2).

Monotonicity makes this a nonincreasing sequence, and equation (III.70.3) supplies one strict bound for every $d \geq 4$. The positive first-variation recurrence proves an obstruction to cosine extremality at each order. It gives neither the rate of decay of c_d^{odd} nor even a positive lower bound uniform in d . Those are distinct extremal questions.

Research question V.113.11 (Stability near the true norm extremizers). If a function has ratio close to c_d^{odd} , what structural information can be forced about its Fourier spectrum?

The dominant-pair parameter p already gives two complementary inequalities, but closeness to our upper estimate should not be confused with closeness to the unknown optimum. A stability theory should begin with an actual extremizer or a proved limiting variational problem. This is also distinct from the near-maximal absolute Gowers-norm problem studied by Eisner and Tao [7].

V.113.3 Returning to the full Gowers argument

Research question V.113.12 (Propagate the local improvements through Section 13). Can the larger spectrum threshold and retained fraction improve the bilinear extraction parameters while maintaining a common grid and the required relation domains?

The first milestone should be a completely stated replacement for one downstream proposition, with its grid, characteristic, and domain hypotheses visible. The original-spectrum corollary offers a conservative route when an existing interface fixes the old threshold. A more ambitious route uses the smaller new spectrum throughout.

Research question V.113.13 (A quantitative repair of the multilinear closure). Which altered induction invariant closes the outstanding Section 16 comparison with synchronized proper partitions?

The repository records that one proposed closing inequality has the wrong direction. Research should therefore seek a new argument or a repaired target bound, rather than attempt to prove that inequality. Candidate invariants should track the graph count, minimum width, common refinement, and integer representatives together.

Research question V.113.14 (An end-to-end explicit density increment). After the remaining phase-removal and partition steps are supplied, what final Szemerédi estimate follows from the strengthened local parameters?

This requires a complete recurrence with all density losses, model changes, rounding terms, and stopping thresholds. The local volume exponent $\alpha^{-3} \log(1/\alpha)$ is useful input, but its effect on the final recurrence cannot be read off from that exponent alone. A strong outcome would be an independently checkable numerical recurrence accompanied by a proof of every transition.

Research question V.113.15 (Comparison with modern inverse-theorem interfaces). Can the original-domain spectral information in the present theorem be retained while substituting a modern quantitative inverse theorem for later parts of the argument?

The quasipolynomial inverse-theorem program [3] is relevant context, but its input and output data differ from this local approximate-homomorphism interface. The research question is to identify a precise bridge and evaluate its full parameter cost. A comparison should use matching hypotheses and outputs before comparing asymptotic labels.

V.114 Source provenance and attribution

Source 28, delivered Section 12.

The comparison baseline is the repository snapshot

`e3125e6982abe2d44bd7cb9369063d4e69d41a2d`.

The source audit included the Section 10 transcription and status ledger, the combined earlier article, the supporting material for later reports, and the incoming characteristic-five, degree-stratification, and sharp threshold manuscripts. The two historical manuscripts most directly relevant to the present result were recovered separately: source 07 for transport-sensitive rigidity and source 12 for dense extension and Bohr radii. This matters because the combined article and its README describe the first five reports and do not exhaust the later collection.

The following attribution boundaries are intentional.

- Gowers’s Section 10 supplies the original local-model strategy, the approximate-homomorphism framework, and the original-spectrum requirement [1].
- Source 07 supplies the prior transport and collision-to-mode viewpoint, including longer paths [133].
- Source 12 supplies prior dense Freiman extension, elementary Bohr regularization, full-fourth-moment smoothing, and the α^5 retention already available in the previous proof [134].
- Source 04 supplies the sharp third-order real norm comparison and the exact quartic cube-count framework [132].

- The present contribution is the logarithmic selection and weighted boundary argument, its directly normalized repair and complete improved parameter chain, and the higher-order norm results proved here.

[write] **6 October 2026; sources 28–33.** *A pin as printed.* The pin of source 07 given in this source’s bibliography ([133]) and in its `data/source_provenance.json` (shipped as `data/28-log-selection-source_provenance.json`) has 39 hexadecimal digits; source 07’s pin is `ae0e895748209f0dd079c6b5ca6df6c6fc065a81`, with a 6 after `6df6c` that is missing here. Recorded, not edited; the other identifiers of that record match the repository (checked at intake). “The combined article and its README describe the first five reports”: at the pin they did; the fifth write of this report covers sources 01–27.

No statement in this report is labelled a resolution of Szemerédi’s theorem’s quantitative optimum. Nor is the cosine calculation described as a refutation of a published conjecture. It settles the concrete question of whether the earlier single-character equality model extends locally to higher orders under the stated hypotheses.

V.115 How to use these results in the ProveIt project

Source 29, delivered Section 10.

V.115.1 Three interfaces, with different domains

The most useful integration is modular. The moment theorem supplies a scalar norm conversion, the stability theorem supplies a structured model together with its selector, and the restriction theorem supplies a large collection of almost entirely respected configurations. Each conclusion has an explicit normalization and domain.

Original location	Present interface	Hypotheses to preserve
Section 3 and Lemma 3.10	Theorems III.68.1 and III.68.3; equation (III.69.7)	The uniform improvement uses real functions with no self-inverse Fourier support. On odd-order groups, centering is sufficient.
Proposition 6.1	Theorems II.110.1, II.111.1, and II.112.1	The hypothesis is proximity to the absolute derivative-correlation endpoint. Selectors have row sums at most one.
Lemma 9.3	The restriction theorems of Section II.120	The ambient domain and codomain are vector spaces over a fixed finite field; all prescribed relations have coefficient sum zero.

Table V.115.1: Connections to Gowers’s paper [1]. The new statements are intended as separately named lemmas with explicit hypotheses.

For the cube application, retain the exact quartic coefficient P_d , then use the smaller factor B_d only when converting $\|f\|_{U^2}^4$ to $\|f\|_{U^d}^4$. This preserves higher-degree terms for whatever counting estimate is being used. In particular, one cannot replace a full cube-expansion remainder by its quartic term without an additional estimate.

For Proposition 6.1, the normalization can be checked directly. An unnormalized derivative sum has N times the normalized Fourier coefficient, and summing its squared magnitude over all h has N^3 times the normalized graph energy. If $|f| \leq 1$ and the full graph sum is at least $(1 - \varepsilon)N^3$, then

$T_\phi(f) \geq 1 - \varepsilon$ and $m = \mathbb{E}|f|^2 \leq 1$. Since $T_\phi(f) \leq m^2$, the relative endpoint deficit satisfies

$$1 - \frac{T_\phi(f)}{m^2} \leq \varepsilon.$$

Thus the joint theorem applies in its stated small-defect range. The supplied graph must be defined on all shifts, or extended in a way whose missing shifts are included in the selector defect.

The repository’s formal statement named `lemma_9_3` is a statement for $\mathbb{Z}/N\mathbb{Z}$ with arbitrary nonzero modulus N . The present restriction theorem is for $\mathbb{F}_q$ -vector spaces. It gives a stronger fixed-field local interface, including polynomial retention and a much smaller threshold, but its field-dependent constant is not uniform in N . For $V = \mathbb{F}_p$, allowing p to grow also makes $S_{p,s}$ grow. In that situation the displayed threshold may be larger than every possible $|B|$. The threshold-free conclusion is still valid, but can be as weak as a singleton. This explains exactly why the finite-field result alone cannot replace the all-moduli lemma inside the full original proof.

[write] **6 October 2026; sources 28–33.** *Checked at HEAD 5d4718772.* `lemma_9_3` is at `Sections08_09.lean`, line 77, a statement over $\mathbb{Z}/N\mathbb{Z}$ with $N \neq 0$, with the exact companion `lemma_9_3_holds` (`Proofs09Restriction.lean`, line 376); `proposition_6_1` is at `Sections06_07.lean`, line 37 (companion `proposition_6_1_holds`, `Proofs05_10.lean`, line 505); `lemma_3_10` at `Sections01_03.lean`, line 340 (companion at `Proofs03CubeUpper.lean`, line 82). The comparison above is accurate; no statement of this source is formalized.

V.115.2 A proposed formalization order

The proofs separate into components that can be formalized and checked independently. A productive order is the following.

1. **Positive Fourier moments.** Define inverse-pair orbits on the non-self-inverse part of the dual group. Prove the balanced multinomial sub-sum inequality with nonnegative coefficients, then derive the Gowers comparison by the existing recursive norm identity. State the torsion profile separately so that no implicit odd-order hypothesis enters the general theorem.
2. **Repair and integration.** Formalize the majority-repair lemma for arbitrary finite abelian groups and the inequality $\delta \geq 3\theta - 6\theta^2$. Next prove the averaged symmetry identity and integrate a symmetric bicharacter using cyclic coordinates. The latter step should use circle-valued phases; division by two inside the coefficient ring would lose the even-order cases.
3. **Projection and endpoint stability.** Prove the two exact Hilbert-space identities for derivative distance and mixed-array distance. The quantitative self, soft, and mixed theorems then follow from norm inequalities, finite Fourier orthogonality, and scalar estimates with rational constants.
4. **Rank and restriction.** Define the full affine-relation space K_x and the error map on that full space. Prove graph rank $= r + e$, the weighted rank-counting lemma, and the expected signed certificate. Keep the prescribed relation space $\mathcal{R}$ distinct from K_x , since they coincide only on generic tuples.

The present package supplies mathematical statements, full proofs, and executable finite diagnostics. It adds no Lean declarations and does not alter the repository’s proof-status ledger. In particular, the experiments below do not certify a new formal theorem.

V.116 Verification and reproducibility

Source 29, delivered Section 11.

V.116.1 What was checked

The three proof families were reviewed independently during preparation. Separate scripts exercise the identities, boundary cases, and examples most vulnerable to normalization or sign errors. Integer and rational checks are distinguished from floating-point diagnostics.

Check	Arithmetic	Coverage
Norm comparisons	Exact rational	65 direct cube-product cases; 480 further Fourier-convolution profile checks; 25 sharp moment cases; 50 odd-cyclic lower-bound examples
Displayed coefficients	Exact integers and rationals	The 222 balanced four-cube choices, the complete three-cube profile, and 50 instances of the interval lower-bound formula
Restriction	Exact finite fields	Five cases, with 294,881 map-tuple incidence checks, including repetitions, degenerate tuples, multiple relations, and vector-valued images
Quadratic integration	Exact rational phases	96 symmetric bicharacters on eight groups; 11,376 polarization identities and 11,376 derivative identities
Majority repair	Exact finite enumeration	All 65,536 maps $\mathbb{F}_2^4 \rightarrow \mathbb{F}_2$, including 240 nonhomomorphisms meeting the repair hypothesis; all 81 maps $C_4 \rightarrow C_3$
Mixed identities and stability	Complex floating point	80 shift/modulation checks, 96 averaged-symmetry checks, 56 self and 56 mixed reconstructions, and six sharp-rate examples

Table V.116.1: Finite diagnostics included with the source. Counts refer to the recorded runs, not to the number of cases in the general proofs.

The direct norm checks multiply the cube-vertex values, providing a calculation independent of the derivative/Jensen route used in the proof. The torsion-profile checks include pure two-torsion, mixed torsion, and odd-order groups. The exact sharpness examples use a character of sufficiently large order; its order need not be prime. The additional odd-cyclic examples have one zero value and equal numbers of 1 and -1 , giving a second realization of the dimension-independent $1/3$ lower limit.

For restriction, the program enumerates the random maps and checks the exact survival probability λ^{r+e} for individual tuples. It also checks both weighted product bounds. The recorded field computations use $\mathbb{F}_2$ and $\mathbb{F}_3$. The proof over arbitrary prime-power fields uses only finite-dimensional linear algebra and the cardinality of affine spans; no extrapolation from these two fields is needed for that proof.

For quadratic rigidity, rational phase arithmetic verifies integration even when the needed phases have order larger than the exponent of the underlying group. The mixed diagnostics verify the signs in the shift/modulation identity and the affine selector $\rho\psi(h)$. Floating-point agreement is only an implementation check of the displayed identities. The output records its observed maximum residual and the separate exact checks.

V.116.2 Contents and build commands

The source package contains the complete L^AT_EX manuscript, its compiled PDF, the vector figure used in the article, a figure generation script, the four verification scripts, their JSON outputs, a source-audit manifest, a short README, and a Makefile. It contains no copies of third-party papers.

From the package directory, the ordinary build command is

```
latexmk -pdf -interaction=nonstopmode -halt-on-error \  
gowers_moments_stability_restriction.tex
```

The verification commands are

```
python code/verify_moments.py --output data/moment_verification.json  
python code/verify_coefficients.py --output data/coefficients.json  
python code/verify_rank_restriction.py --output data/rank_verification.json  
python code/verify_joint_rigidity.py --output data/joint_verification.json
```

The moment, coefficient, and restriction checks use Python’s standard library. The joint numerical diagnostics require NumPy. Regenerating the scientific figure requires Matplotlib; the supplied figure is already sufficient to compile the manuscript. The Makefile provides `all`, `verify`, `figures`, and `clean` targets.

V.117 Further research questions

Source 29, delivered Section 12.

The questions below are unresolved by the present arguments and, where specified, by the inspected project reports. They are proposed research targets rather than assertions of exhaustive priority or universal open-problem status.

V.117.1 Sharp higher-norm comparisons

Question V.117.1 (Determine the fourth-dimensional constant). *What is the exact value of c_4^{odd} ? The present bracket is*

$$(8/111)^{1/4} \leq c_4^{\text{odd}} \leq (8/35)^{1/4}.$$

Is a single Fourier inverse pair extremal, or can a more distributed spectrum produce a larger ratio?

[write] **7 October 2026; sources 60–69.** The exact value at $d = 4$ is still open after sources 60–69; the bracket is $0.5587725 \leq c_4^{\text{odd}} \leq (4/29)^{1/4} = 0.6094183$ (item III.1), and source 63’s limit theorem does not change it.

The autocorrelation moment theorem is sharp, but its equality model has a substantial gap in the recursive norm estimate. A sharp proof must control these two losses together. A useful first target is an improved inequality for $\|\Delta_h f\|_{U^3}^8 - |c_f(h)|^8$ averaged over h , expressed in Fourier quantities

compatible with the pair-concentration estimate. The direct four-cube expansion offers a second route, through positive families of balanced vertex labelings.

Question V.117.2 (The limit of the extremal constants). *Determine*

$$c_{\infty}^{\text{odd}} = \lim_{d \rightarrow \infty} c_d^{\text{odd}}, \quad \frac{1}{3} \leq c_{\infty}^{\text{odd}} \leq \frac{1}{2}.$$

Can the upper bound $1/2$ be approached by a family of real centered functions as the dimension increases?

[write] **7 October 2026; sources 60–69.** *Answered by source 63:* $c_{\infty}^{\text{odd}} = 1/3$ (its Theorem III.97.1); source 60 had excluded $1/2$ (Section III.94).

The proof of the upper bound takes a supremum over groups and functions for each d . One cannot interchange that supremum with a limit for a fixed function without justification. A productive experiment would compare long interval models, sparse Fourier models, and bounded functions with increasingly complicated sign patterns. Any numerical search should report the group, the dimension, and the normalization separately.

Question V.117.3 (An optimal profile at fixed self-inverse mass). *For a prescribed τ , determine the sharp bound for $\|f\|_{U^2}^4 / \|f\|_{U^d}^4$ among real functions having self-inverse Fourier fourth-mass fraction τ .*

The polynomial $K_m(\tau)$ solves the corresponding uniform autocorrelation-moment problem. Passing to the Gowers norm can introduce a further gap, just as it does at $\tau = 0$. It would be useful to know whether a sharp norm profile depends only on τ , or whether finer information about the orders and relations of the active characters must also enter an effective bound.

Question V.117.4 (Stability for actual norm near-minimizers). *Can a function whose ratio is close to the sharp higher-norm constant be classified up to a controlled error in a natural norm?*

Proposition III.68.2 gives U^2 concentration only when the autocorrelation moment is near its own lower bound. It is not a stability theorem at the unknown sharp Gowers ratio. A successful result would identify the relevant models first and then show that near equality in all necessary steps forces proximity to those models. For an L^2 conclusion, additional amplitude or spectral-tail information may be needed.

V.117.2 Joint rigidity beyond the present endpoint

Question V.117.5 (Optimal constants and a larger defect range). *What are the best constants in the joint function-selector stability theorem, and how far can its small-defect range be enlarged while retaining a single global quadratic model?*

The 24 and 192 constants are explicit common bounds, not optimized values. The positive-amplitude example forces squared function error of asymptotic order $\varepsilon/2$; the soft-selector example forces missing graph mass of order ε . Sharper homomorphism-testing theorems, such as [122], can refine the repair step. The remaining losses come from the derivative projection and the passage

from the repaired graph to one Fourier atom. Studying the exact two-point and cyclic examples may indicate which losses are intrinsic.

Question V.117.6 (Remove the row constraint in a natural normalization). *Find a joint stability theorem for ambiguity or mixed short-time Fourier transforms when the selector is not required to have row sums at most one.*

The class of exact models must then change with the normalization. Subgroup-localized quadratic phases and isotropic phase-space supports occur in the existing exact theory [123]; critical- L^p Gowers near-extremizers likewise allow coset localization [7]. A well-posed next theorem should first specify its energy normalization and model class, then measure distance both in function space and in phase-space support. Extending the current statement verbatim would suppress genuine exact examples.

Question V.117.7 (Higher-degree joint derivative selectors). *For $d \geq 4$, can near-maximal capture of $(d - 1)$ -degree derivative structure recover a degree- d phase together with the supplied family of derivative models, with a linear squared-error bound?*

The quadratic proof repairs a map into an abelian character group and then integrates a symmetric bicharacter. Higher degree introduces compatibility among several differences and, in small characteristic, nonclassical circle-valued polynomial phases. A useful first case is a finite vector space over a field of characteristic greater than the degree, with all derivative models supplied on the full domain. The challenge is simultaneous consistency of the models, beyond pointwise applications of an inverse theorem.

Question V.117.8 (Effective reconstruction and certification). *Can the joint theorem be converted into an algorithm that reconstructs Q, ψ, ρ from approximate derivative data and outputs a checkable error certificate?*

The proof already identifies operations: choose large Fourier coefficients, apply majority repair, test symmetry, integrate the bicharacter, and find a Fourier atom after twisting. Exhaustively performing those operations is finite but can be expensive. The algorithmic questions concern sample complexity, treatment of noisy coefficients, and time bounds in the group size or its presentation. For a small fixed group, an exact reconstruction tool would also be a useful bridge to formal verification.

V.117.3 Sharper and broader restriction

Question V.117.9 (Optimize the complete rank profile). *Given the counts $G_{r,e}$ and $U_{r,e}$, optimize the exact certificate*

$$\sum_{r,e} (\eta G_{r,e} - U_{r,e}) \lambda^{r+e}$$

over the admissible values $\lambda = q^{-t}$. What information less expensive than the full profile still gives a substantial improvement over the worst-case coefficient?

The present theorem replaces all respected graph ranks by s and all unrespected error ranks by one. Data with few low-rank bad tuples or many low-rank good tuples can behave much better. The signed polynomial is already an exact certificate, so concrete improvements can be proved without changing the random selector. The last coordinate of that selector already supplies a uniform affine offset. Moreover, a mixture of restriction depths gives a convex combination of the corresponding

signed certificates, so it cannot improve their maximum for a fixed full profile. A remaining question is whether a different selection family can avoid the loss associated with discrete depths.

Question V.117.10 (Quantify degeneracy more geometrically). *Can the product $z \prod_{i=0}^{s-2} (z + q^i)$ be replaced by a bound in terms of actual affine-span intersection statistics of B , or by a sharper universal polynomial?*

At each step the proof permits up to q^{j-2} points in the current affine span, independently of earlier choices. This is sufficient for the linear threshold but often very wasteful. Cap sets, subspace-evasive sets, and sets with small intersections with low-dimensional flats should give stronger estimates. Any improvement for the generic higher-error-rank theorem must still control low-domain-rank unrespected tuples; a generic rank gap alone says nothing about them.

Question V.117.11 (An efficient deterministic restriction). *Can one find a map A meeting the expected signed certificate in time polynomial in a useful input representation?*

Exhaustive search over all linear maps proves effective existence. Conditional expectations may derandomize the proof, but computing the conditional potential from the full tuple list can already cost $|B|^\ell$. A meaningful algorithm should exploit the relation matrix, Fourier structure, or a compressed rank profile. It should also distinguish the task of exhibiting B' from the task of certifying its proportion of respected relations.

Question V.117.12 (A useful growing-modulus analogue). *Can actual-rank or multiplicity-sensitive survival improve the prime-cyclic restriction theorem with constants that remain useful as the prime grows, and eventually handle arbitrary cyclic moduli?*

This is the principal unresolved obstacle to using the fixed-field improvement as a replacement for the original Lemma 9.3. The prior manuscript [136] also contains a Fejér-filter method in cyclic groups. A possible direction is to track the survival of repeated or low-rank configurations inside that filter, while accounting for approximate relations and torsion. The finite-field product proof by itself gives no such uniformity.

V.117.4 From local estimates to a global bound

Question V.117.13 (A complete quantitative substitution). *Which of the local improvements in the accumulated project reports survive a full quantitative inverse and density-increment argument, and what explicit progression-free-set bound follows?*

This question requires a dependency calculation, not just a list of smaller constants. Every invocation must preserve its domain, normalization, density loss, length threshold, and error budget. The fixed-field and cyclic branches should be analyzed separately. A useful deliverable would be an executable symbolic recurrence, verified against the corrected statement of each invoked lemma, and then compared with an established global result such as [2]. The present article supplies several improved local inputs for that project.

Among these directions, three immediate targets have especially clear interfaces: improving the explicit c_4^{odd} bracket, optimizing the complete finite-field rank certificate on structured sets, and formalizing the joint stability theorem over arbitrary finite abelian groups. Each can be pursued independently and tested against a fully specified theorem above.

V.118 Repository crosswalk and exact validation

Source 30, delivered Section 10.

V.118.1 Pinned interfaces

The repository pin used for the direct statement comparison is

8d65d3c9febe4d0e9efe945e11b53af5f19d1435.

The relevant files are under `Combinatorics/Ramsey/Lean/GowersSzemerédi/`. The inspected blob identifiers are

File	Git blob identifier
<code>Definitions.lean</code>	97113b7afa6925a2dd4b76641eeaeff09597ab6a
<code>Section10.lean</code>	302e8a223f56dcdabf29f30ca3c80ac62a7adbfc

The definition of `bohr` uses the centered absolute residue of kd bounded by ρN , exactly matching the normalization here. `section10Translate` translates by addition of d . `InducesDifferenceMap` expresses the identity $\psi(x) - \psi(y) = \psi_1(x - y)$ used in this manuscript.

Repository interface	Proposed mathematical companion
<code>lemma_10_10</code>	Theorem II.130.2, first for cyclic groups and $0 < \rho \leq 1/4$; retain the intrinsic shell inequality as a reusable lemma.
<code>corollary_10_11</code>	Corollary II.132.1 with unchanged overlap margins; Corollary II.132.2 as a separate larger-domain theorem.
<code>lemma_10_12</code>	Theorem II.133.1 and Corollary II.133.2, with all finite-fibre nonemptiness hypotheses explicit.
<code>theorem_10_13</code>	A future propagation target, not a theorem proved by this package. Its spectrum and parameter construction must be audited separately.

The earlier local-quantitative-refinements README and the transport-rigidity integration note were inspected to avoid simply repeating the existing work on phase partitions, cube expansions, transport rigidity, and restriction arguments. That inspection is not an exhaustive text comparison with every pending manuscript. The source audit accompanying this package distinguishes pinned reads from earlier moving-branch reads.

V.118.2 Formalization route, not claimed formalization

The shortest plausible implementation begins with finite sets in $\mathbb{Z}/N\mathbb{Z}$, not with a general Haar-measure library. Useful stages are:

1. Prove the two exit-cardinality identities and the abstract disjoint translation lemma. It is best to index the $2m$ tubes by a sign and a positive integer, so that same-sign and opposite-sign collisions are separate proof cases.

2. Establish a real or integer centered-residue lift lemma through $2m$ steps. This is the only non-return ingredient and is where the $\rho \leq 1/4$ hypothesis is used.
3. Prove the three-cell partition with unambiguous endpoint assignments. Translate each nonempty cell to inject it into the original Bohr set. Do not count only the image of the character map without accounting for its fibres.
4. Combine the cardinal inequalities before introducing real division. Keep the exact intrinsic form and obtain the 3^r form as a corollary.
5. Prove the general tree budget and the induced difference identity; then translate the tuple formulation of Freiman relations into the repository's exact API.
6. Prove the compatibility theorem with the two exceptional direction sets measured against $|A|$. Keep positivity of the common fibre as an explicit premise.

No Lean file has been compiled for the new results in this package. The presence of a proposition-valued declaration in the existing catalogue is not a substitute for a proof. The proposed theorem names and dependency order in `FORMALIZATION.md` are implementation suggestions, not a report of successful kernel checking.

[write] **6 October 2026; sources 28–33.** *Checked at HEAD 5d4718772.* `Definitions.lean` and `Section10.lean` have the blobs recorded above; `lemma_10_10`, `corollary_10_11`, `lemma_10_12` and `theorem_10_13` are at lines 362, 374, 401 and 473 of `Section10.lean`, with exact companions `lemma_10_10_holds` and `corollary_10_11_holds` (`Proofs10Bohr.lean`, lines 277, 607), `lemma_10_12_holds` (`Proofs10Induced.lean`, line 145) and `theorem_10_13_holds` (`Proofs10Main.lean`, line 1292); `bohr` is at `Definitions.lean`, line 282. No statement of this source is formalized; the names in its `FORMALIZATION.md` (shipped as `30-fixed-radius-bohr-FORMALIZATION.md`) are proposals. The recorded output `results/verification.json` named above is shipped as `data/30-fixed-radius-bohr-verification.json`.

V.118.3 Exact finite checks

The delivered `code/verify.py` uses only the Python standard library, integers, and rational fractions. It uses explicit exceptions rather than `assert`, so optimization mode does not disable validation. The run recorded in `results/verification.json` completed the following checks:

Suite	Recorded coverage
Cyclic Bohr sets	13 404 parameterized sets; 135 182 admissible shift/length instances, of which 11 642 have positive exit count.
Unequal-width product groups	500 seeded parameterized sets; 4 588 admissible instances, including 676 with positive exit count.
Four-vertex budgets	269 759 finite intersection inequalities, including repeated positions and negative lower bounds.
Sharpness	Exact interval equality and exact 3^r packing examples.
Scalar and local-map tests	1 140 rational parameter checks; a nonconstant map from 73 points into $\mathbb{Z}$, with induced orders two and three on a nine-point neighborhood.
Spiked sets	Exhaustive product-group membership for $s = 1, 2, 3$; explicit cyclic parameter certificates for $s = 1, \dots, 6$.
Cyclic brute force	Direct enumeration of the rank-three CRT example with $N = 90\,799$, $ B = 19$, and defect 3.

The cyclic suite enumerates moduli 2 through 32, with character representatives from 0 through $\lfloor N/2 \rfloor$, six rational widths, and $m \in \{1, 2, 3, 4, 8\}$. It includes lists of size at most two throughout and lists of size three through modulus 17. Opposite frequencies give the same single-coordinate constraints, which is why these representatives suffice for the stated suite. The product suite is seeded, not exhaustive. Every qualifying tube is explicitly checked for non-return, pairwise disjointness, and containment in the *direction-dependent* enlarged shell.

The large CRT certificates check all elementary parameter inequalities, pairwise gcd conditions, distinct explicit frequencies, and exact step values. They do not enumerate their enormous ambient groups; the general membership and cardinality formulas for those groups are established by Theorem II.135.3. The finite checks are diagnostics and reproducibility evidence, not proof-assistant certificates or an independent review of the general argument.

V.119 Research questions and next integration targets

Source 30, delivered Section 11.

The first three questions concern the central quantitative limitation. The later ones concern effective use and formal integration. None is assumed in any proof above.

Question V.119.1 (Optimal exponential base). *Determine the asymptotic growth of $\mathcal{L}_r$. Is its exponential rate $\sqrt{3}$, 3, or another value? Even improving either side of (II.135.8) by an exponential factor would clarify which part of the geometry the box-packing argument loses.*

Question V.119.2 (Prime cyclic groups). *What is the analogue of $\mathcal{L}_r$ when the ambient group is restricted to $\mathbb{Z}/p\mathbb{Z}$ with p prime? Can one realize an exponential boundary obstruction without the transverse Chinese-remainder factors, or does prime order permit a substantially better universal bound?*

Question V.119.3 (Joint sharpness of escape and packing). *Classify, even approximately, the Bohr sets for which $2mb_B(d)$ is close to $|B(3\rho/2)| - |B|$ and the latter is close to $(3^r - 1)|B|$ simultaneously. The separate sharpness examples show why optimizing the two inequalities independently is insufficient.*

Question V.119.4 (Intrinsic growth from spectral structure). *The spectra produced in Section 10 are not arbitrary lists of characters. Can their additive structure, or the structure of the fibre-count function that generates them, force D_B far below 3^r at the original radius? A positive result could combine the intrinsic theorem with information already present upstream, rather than use worst-case rank alone.*

Question V.119.5 (Adaptive anisotropic widths). *When some characters constrain the proof much more tightly than others, can one choose unequal widths before the partial map is constructed so as to optimize the final extension neighborhood? A correct optimization must preserve every earlier invariance premise; choosing widths after the map has been constructed can change the problem.*

Question V.119.6 (Optimal dense-domain extension budget). *For exact Freiman maps, how far can the sufficient condition $(2k - 1)\kappa + 2k\eta < 1$ be improved? Can overlaps among the exceptional sets be exploited uniformly, or are there examples where the tree budget is close to optimal? The question should separate failure of set inclusion $C \subseteq A - A$ from failure of the higher-order relations.*

Question V.119.7 (Compatibility thresholds). *Is the $5/14$ threshold in Corollary II.133.2 optimal under only the explicitly stated pair-overlap and finite-fibre hypotheses? Can target-fibre information*

improve the condition $\theta < 1/2$ when the two good-value sets are not arbitrary subsets of the common fibre?

Question V.119.8 (Propagation through Theorem 10.13). *Re-run the complete corrected proof with a separate radius parameter provided by Corollary II.132.1. Determine an explicit replacement for the final local-model radius, while retaining the corrected integer spectrum bound, density conclusion, and nonempty-fibre conditions. This is a concrete next theorem, but it has not been established by a symbolic substitution in this manuscript.*

Question V.119.9 (Global relevance). *After such propagation, which losses dominate the final quantitative Szemerédi theorem? Does removing this power-of-radius loss change only constants, an iterated exponent, or a genuine asymptotic class in a fully specified implementation of the argument? This requires an end-to-end parameter ledger, not just a local comparison table.*

Question V.119.10 (Beyond rectangular character constraints). *For which other finite sets does a small translation admit the same $2m$ non-return property? Can convex constraints on several character coordinates be handled using directional widths, with a covering number replacing 3^r ? The abstract escape lemma already applies once the non-return and enclosure properties are proved.*

Question V.119.11 (Efficient exact growth certificates). *Can one certify an upper bound for D_B substantially below 3^r without enumerating all of G ? Lattice descriptions or compressed descriptions of the character-map fibres may provide smaller exact certificates for large cyclic moduli.*

Question V.119.12 (Kernel-checked integration). *Implement the intrinsic escape theorem, packing theorem, and margin-preserving extension in the repository’s toolchain, followed by an axiom audit. The most important proof obligations are endpoint-safe non-return and opposite-sign tube disjointness; the final scalar inequalities are comparatively small.*

V.120 Conclusion

Source 30, delivered Section 12.

The decisive change is to count translated exits inside a controlled Bohr enlargement, rather than count a boundary in the ambient group and then divide by a small density. The resulting inequality retains the original set and replaces a power-of-width loss by an exponential-rank loss. It gives a substantially larger neighborhood for the dense-domain Freiman step, supports higher-order extension and exact compatibility budgets, and has a matching qualitative obstruction: exponential rank cannot in general be replaced by polynomial rank without further hypotheses. The exact exponential rate and the prime-order case remain specific, well-defined research targets.

V.121 Build, verification, and file layout

Source 30, delivered Appendix A.

The manuscript is self-contained as a single `article.tex`; its bibliography is embedded below. No source paper, repository snapshot, or font file is redistributed. A standard TeX Live or MiKTeX installation with the packages named in the preamble is sufficient.

```
python3 code/verify.py --output results/verification.json
latexmk -pdf -interaction=nonstopmode -halt-on-error article.tex
```

The delivered package also contains `README.md`, `SOURCE_AUDIT.md`, `FORMALIZATION.md`, `Makefile`, the exact verification script, and its recorded output. The Python check suite is deterministic; the product-group sample uses seed 20261006. Build logs and render checks are summarized in `results/build_validation.json`.

V.122 Verification and a proposed formal integration

Source 31, delivered Section 10.

The mathematical proofs are the justification for the universal claims. The accompanying programs check finite identities, constructions, normalizations, and endpoint constants. They distinguish exact integer or rational checks from complex floating-point diagnostics. None of the programs is a proof assistant.

Program	Purpose
<code>code/verify_spectral_rigidity.py</code>	Weyl mass and self-duality; subgroup and phase examples; the quartic equality and inverse formulas
<code>code/verify_weighted_extraction.py</code>	Finite weighted extraction; vertical-defect injection; prime-field progression selection and the order- k condition
<code>code/verify_weight_defects.py</code>	Exact defect inequality and counterexample formulas; selected-graph endpoint diagnostics

The recorded JSON files in `data/` state the actual test ranges and numerical tolerances. Run `makeverify` to reproduce them. The article compiles with pdfLaTeX through `makepdf`, without shell escape or a network dependency. A source ledger records the repository revision and the provenance of the principal inputs.

V.122.1 Interfaces already present in the repository

At the inspected revision, `Sections06_07.lean` defines `proposition_6_1`, `lemma_7_5`, and `corollary_7_6` as propositions. The file `Proofs07FreimanRestriction.lean` proves `lemma_7_5_holds`; its private vertical-fiber bound uses C^{4k} . The proposed change to that local proof is concrete: use (II.116.3), then use the finite progression selection of Theorem II.116.1, and retain weights in the averaging step. Existing statements can be recovered as weaker corollaries.

[write] **6 October 2026; sources 28–33.** *Checked at HEAD 5d4718772.* `proposition_6_1`, `lemma_7_5` and `corollary_7_6` are at lines 37, 183 and 197 of `Sections06_07.lean`; `lemma_7_5_holds` is at `Proofs07FreimanRestriction.lean`, line 267, and its private bound `cor75_verticalFiber_bound` (line 110) is C^{4k} , as the text says. The formal `lemma_7_5` retains $|B|/(8kC^{4k})$ for a prime modulus; Theorem II.116.1 implies its conclusion, as source 10’s theorem already does. Nothing here is formalized.

The spectral argument needs a separate finite complex inner-product layer: Weyl operators, symplectic orthogonality, common eigenspaces, and finite character projectors. The tail-energy identities should be proved first over nonnegative finitely supported functions on an abelian group. This isolates the exact polynomial algebra from the linear-algebra proof of the extremal subgroup.

V.122.2 Suggested order of implementation

1. Formalize the weighted graph extraction and the actual vertical defect set. These mostly use finite sums, finite sets, and the Plünnecke–Ruzsa interface already needed by the repository.
2. Prove (II.107.6), the unequal-pair estimate for C_0 , and the quartic deficit identity over an abstract finite abelian label group. Keep every subtraction in the reals or in a division-free nonnegative formulation.
3. Establish the Weyl mass and self-duality identities, followed by the $3/2$ small-difference lemma and the noncommuting subgroup gap.
4. Prove the trace formula for character projectors and identify their one-dimensional ranges with quadratic coset phases.
5. Assemble the sharp spectral theorem; prove monotonicity and the two quadratic inversions separately. Then derive the full and partial graph corollaries.
6. Formalize the weight-defect decomposition and the two finite counterexample families. Complexity-dependent L^1 estimates and measurement transfer can be added afterward.

This is an integration plan, not delivered Lean code. We did not run the repository’s Lean build and do not alter its proof-status ledger. An ordinary proof, a correct match to the source statement, and kernel verification are distinct requirements.

V.123 Further research questions

Source 31, delivered Section 11.

Research question V.123.1 (The full critical-norm profile). Determine the optimal universal function relating $\kappa(f)$ and $F(f)$ beyond $1 - \kappa < 1/120$. Does the two-state curve remain extremal on a larger interval, and at which first value do other configurations take over? The cutoff here is sufficient, not a candidate optimal transition.

Research question V.123.2 (Odd-order sharpness). Find the optimal second and higher coefficients in the odd-order infidelity profile. Theorem II.107.3 excludes the universal order-two equality mechanism, but its odd-order envelope is not claimed sharp. A two-state configuration of odd-order label has a different outside fourth moment, suggesting a more refined analysis of phase interference.

Research question V.123.3 (Stability of the extremizers of the envelope). Turn (II.107.4) into a quantitative classification of all near-equality states by their distance to an embedded two-state configuration. The tail concentration is explicit; the remaining task is a sharp phase-alignment estimate for the outside Weyl coefficients.

Research question V.123.4 (The optimal joint graph budget). The separate graph and phase bounds in (II.108.2) are sharp. Determine the exact feasible region of $(\varepsilon, t, z, \beta)$. In particular, replace the valid but not certified-sharp coefficient $1 - 8t(1 - t)$ in (II.108.1) by its optimal counterpart, or identify the additional statistic needed for such a formula.

Research question V.123.5 (The actual defect-set geometry). In Theorem II.116.1, optimize the weight retention using the set V_k itself rather than only $|V_k|$. Intervals are a convenient way to avoid nonzero defects, but other target sets may retain substantially more weight for structured V_k .

Research question V.123.6 (Torsion-sensitive target restriction). Extend the finite weighted restriction to composite cyclic and general finite abelian targets. The prime proof uses invertibility of the integers $1, \dots, kL$. A general result should charge the sizes of their kernels and the distribution of the defect set across those kernels, rather than silently importing the prime case.

Research question V.123.7 (Better weighted extraction exponents). Use the square-root mass theorem in Appendix II.118, or a sharper weighted BSG theorem, to improve the powers in (II.116.4) while preserving the originally specified weight. Separate the loss from structural extraction, the choice of one point per fiber, and exact Freiman restriction.

Research question V.123.8 (Replacing the worst fiber size). Can K^{2k+2} be replaced by a profile involving the masses and largest atoms of individual fibers? The present proof already admits the actual mass of the maximum-weight graph as an input. A stronger theorem should retain this distribution through the small-doubling step as well.

Research question V.123.9 (Compatible restrictions across slices). Construct the exact Freiman restrictions for different derivative slices compatibly. This is the missing link between the local correlation-preserving theorem and Gowers' later multilinear integration. A favorable exponent for each slice alone does not establish a favorable global exponent.

Research question V.123.10 (Weights under weaker complexity assumptions). Replace finite Fourier support and a finite number of occupied cosets in (II.109.2) by effective entropy or tail norms. The two counterexamples show that some such hypothesis is necessary. A useful formulation should recover L^1 control without demanding literal sparse support.

Research question V.123.11 (The optimal coupled weight defect). Determine the sharp dependence of $\mathcal{E}(w) - a^4$ on the norm deficit δ , the energy-saturation error η , and the mean a . Proposition II.109.4 gives an exact linear transfer identity. The subsequent fourth-power estimate in (II.109.7) need not be jointly optimal; retain the mean and the eigenbasis collision mass to sharpen it.

Research question V.123.12 (Mixed functions and higher norms). Classify the equality and near-equality pairs in the mixed derivative inequality for two different input functions. For higher Gowers norms, identify an analogue of the spectral subgroup construction and the finite tail-energy envelope. Neither extension follows merely by changing the exponent eight.

Research question V.123.13 (A checked quantitative application). Choose one fixed progression length, integrate the new local interfaces into a complete dependency ledger, and verify every resulting threshold. Only after that calculation would it be justified to claim an improved bound produced by this proof route. Comparison with modern bounds would then be a further step.

V.124 A short dictionary for the extremal geometry

Source 31, delivered Appendix B.

Additive-combinatorial description	Finite operator description
Derivative Fourier coefficient	Weyl matrix coefficient
Normalized squared derivative spectrum q/N	Characteristic-function probability distribution of a pure state
Quadratic phase on a coset	Common eigenstate of a maximal commuting Weyl subgroup
Full-support quadratic phase	Extremal subgroup projecting bijectively onto the translation coordinate
Critical ratio κ	Normalized fourth moment of Weyl coefficients
Tail labels relative to an extremizer	Other characters in its common eigenbasis

This dictionary explains the relevance of the stabilizer literature without imposing a prime-power assumption on the finite abelian group. All operator arguments used in the article were proved directly. For general finite abelian groups, “stabilizer” here means the maximal commuting Weyl-eigenstate description; no transitivity claim for a particular definition of a Clifford group is needed.

V.125 The pinned source and the literature

Source 32, delivered Section 6.

V.125.1 The full degree two source implication

The precise repository target is `corollary_5_8` in `Section05.lean`, lines 242–259, at commit

`c94998a5db18ba63c815a3bf6c3877a0fde91930`.

Direct links are in the references and `SOURCES.md`. This is a comparison with the predicate at that pin, not a claim that its general-degree proof has been formalized or that a repository build was rerun.

The source constant is

$$K = (k!)^2 2^{(k+1)^2}, \quad K|_{k=2} = 4 \cdot 512 = 2048.$$

Its degree-two conclusion is a proper progression with length at least $S^{1/2048}/8$ and density at least $\delta + \alpha/16$. The source uses this corrected gain rather than the original paper’s displayed $\alpha/8$.

[write] **6 October 2026; sources 28–33.** *Checked at HEAD 5d4718772.* `Section05.lean` and `Definitions.lean` have the blobs read at the pin (882b17a, 97113b7); `corollary_5_8` still starts at line 242 and is open in the ledger, which records 105 exact companions and 15 open statements. The translations below are accurate: `ModAP.IsProper` (`Definitions.lean`, line 93) is injectivity of the parameterization. The SHA-256 values in `SOURCES.md` (shipped as `32-quadratic-increments-SOURCES.md`) are of the files with one line feed appended, the retrieval convention of source 26’s record; source 41’s placement records the same finding.

The hypotheses translate without an extra restriction:

1. The source’s `PolynomialOn 2 Finset.univ` means there exist three coefficients with $\phi(x) = c_0 + c_1x + c_2x^2$ for every x . This is exactly the coefficient-polynomial convention used here, including composite moduli.

2. The definition `ModAP.IsProper` is equality of carrier cardinality and parameter length. Since the carrier is the image of a finite parameter set of that size, this is exactly injectivity. Despite an informal source comment mentioning wrapping, the definition does not require nonwrapping standard representatives.
3. Empty parents are syntactically permitted by the source, but cannot occur with its other premises. If one parent is empty, two-comparability forces every parent to be empty. Then the cover forces A to be empty, and all balanced correlations vanish, contradicting $\alpha N > 0$.
4. The source assumes $N \neq 0$. Its $N = 1$ case has identically zero balanced function for either possible subset, so positive correlation is impossible. For $N \geq 2$, the identity $|A| = \delta N$ identifies its δ with the density here. Its balanced function and exponential are the same $\mathbf{1}_A - \delta$ and e_N used above.
5. Disjointness, coverage, two-comparability, and the summed correlation agree literally. The source does not require the concluding progression to lie in the cover.

The prefactors 17 and 8 require an explicit comparison; comparing exponents alone is insufficient. In fact, for every $S \geq 1$,

$$\max \left\{ 1, \frac{S^{1/39}}{17} \right\} \geq \frac{S^{1/64}}{8} \geq \frac{S^{1/2048}}{8}. \quad (\text{V.125.1})$$

If $S^{1/64}/8 \leq 1$, the first inequality follows from the singleton bound. Otherwise $S > 8^{64}$ and

$$S^{1/39-1/64} > 8^{25/39} > \sqrt{8} > 17/8.$$

Here $25/39 > 1/2$, and $8 > (17/8)^2$ is the integer inequality $512 > 289$. Dividing gives $S^{1/39}/17 > S^{1/64}/8$. The second inequality in (V.125.1) uses $S \geq 1$. The gain $\alpha/3$ is greater than $\alpha/16$. Thus the very same progression from Theorem F.37.1 satisfies the complete pinned degree-two conclusion at every permitted scale.

V.125.2 Primary literature comparison

Gowers [1], Section 5, develops Weyl-type recurrence and polynomial-to-affine localization; Lemmas 5.3–5.5 and Corollaries 5.6–5.8 are the relevant chain. Corollary 5.8 on printed page 493 concerns the same disjoint comparable-cover mechanism. The present proof uses its own explicit quadratic recurrence and a second-moment alternative, rather than importing a large-threshold polynomial localization statement without verifying its scale premise. The comparison above is specifically to the edited repository predicate [142], whose gain is $\alpha/16$.

Green and Tao [128], Appendix A, Proposition A.2, give simultaneous quadratic recurrence for real $\alpha_1, \dots, \alpha_d$: there is $1 \leq n \leq X$ with $\|n^2 \alpha_j\| \ll dX^{-c/d^2}$. The following discussion records Schmidt's stronger asymptotic exponent $-1/(d^2 + d) + \epsilon$ and explains why dependence of implicit constants on dimension matters. Their Section 6 treats linearization. These results confirm the classical nature of recurrence followed by localization. Our weaker single-coefficient estimate has a fully displayed uniform constant and is proved here to keep the numerical interface self-contained.

Maynard [72], Theorem 1.1 and Corollary 1.2, treats systems of real polynomials with zero constant term. For m polynomials of degree at most d , the corollary gives simultaneous fractional parts $\ll_{d,m} X^{-c_d/m}$ for some $1 \leq n < X$, uniformly in their coefficients. This is a substantially

more general recurrence theorem; the displayed corollary leaves its constants implicit. We neither claim that R^5 is the best quadratic recurrence rate nor infer any optimality for $1/39$ from it.

The specific claim established here is the explicit, all-moduli, threshold-free local-cover theorem and its pinned source implication. Priority for that exact formulation has not been established. No argument here turns this local improvement into a better final Ramsey or Szemerédi estimate, handles all polynomial degrees, or supplies an integer-rectification theorem.

V.126 Reproducibility and further questions

Source 32, delivered Section 7.

V.126.1 What the companion establishes

The accompanying Python program uses integer and rational arithmetic for bounded structural checks. It tests progression injectivity, subgroup-coset pair counts and variance, removal of nonproper directions, cover-mass identities and positive-cell extraction, residue-chain chunking, quadratic phase lifts, reduced curvature, and the reduced-midpoint regression. It also checks the integer identities behind the constants and root-free ceiling choices. Small rational recurrence examples exercise the recurrence interface; they are not a proof about every real coefficient.

The full mathematical proof is in Sections I.45–I.48. The companion does not enumerate a partition of size $2^{116}L^{38}$, even at $L = 2$, and it does not approximate complex exponentials and call the result a certificate. Its finite coverage, argument caps, and limitations are reported in its JSON output and documented in `README.md`. Input-validation failures remain runtime errors under optimized Python. Tests and diagnostics run under both normal Python and `-O`; deterministic diagnostic output must agree.

The offline builder verifies an exact source inventory and, for a packaged distribution, the manifest hashes. It rebuilds the PDF using an installed trusted TeX toolchain with shell escape disabled and writes a deterministic ZIP with fixed member order and metadata. It rejects unexpected source members, source aliases, and unsafe output locations, and bounds subprocess time and output. These are reproducibility and regression controls for trusted source, not a sandbox for hostile TeX or Python and not a guarantee against an actively malicious concurrent filesystem owner. Hashes establish byte agreement with the manifest, not authorship or mathematical truth.

The ZIP contains only this article, its PDF, the original companion and tests, the builder, public documentation, and its manifest. It contains no third-party paper or private review material. The exact commands and platform requirements are in `README.md`; the archive digest is supplied outside the archive to avoid a circular hash.

V.126.2 Questions suggested by the proof

1. **Better recurrence with explicit small-scale constants.** If an unconditional estimate $p \leq CR^\kappa$ is available, the present choices $H \asymp L^3$ and $R \asymp L^7$ make the partition cost proportional to $L^{7\kappa+3}$. The large-correlation cover step loses one further factor of L , suggesting length exponent $1/(7\kappa+4)$ with appropriately checked constants. This is parameter bookkeeping for a future argument; Remark I.45.4 does not supply $\kappa = 4$.

2. **A more economical quadratic partition.** The two-stage construction first suppresses curvature and then suppresses an affine slope. Can these stages be combined to improve the power 38, while retaining every point and allowing arbitrary real coefficients? Any proposed replacement must preserve the small-scale rounding and properness checks.
3. **Parent-contained small-correlation increments.** The variance alternative uses the whole group. Requiring $Q \subseteq U$ or $Q \subseteq P_i$ would be an additional conclusion and needs a separate argument or an obstruction; it is not implicit in the present theorem.
4. **The next polynomial degree.** The source bridge is complete for degree two. Extending the threshold-free strategy to degree three would require an explicit cubic recurrence/localization budget and a fresh check that the large-correlation branch pays for it. Nothing in the degree-two proof establishes that extension.

V.127 Further research questions and a concrete program

Source 33, delivered Section 9.

The questions below are proposed by the present analysis. Unless an external source is specifically cited, they are not being represented as previously published open problems. Each is separated from the proved claims above.

Research question V.127.1 (The optimal common-prime quadratic coefficient). For fixed even s , determine the smallest coefficient c_s^{prime} for which a uniform small-error estimate of the form

$$\delta \leq \varepsilon_s/s + c_s^{\text{prime}}\varepsilon_s^2 + O_s(\varepsilon_s^3)$$

holds on $\mathbb{F}_p^2 \rightarrow \mathbb{F}_p$ for all sufficiently large odd primes. Theorems F.38.2 and II.141.3 give the rigorous bounds

$$\frac{3s-2}{4s^2} \leq c_s^{\text{prime}} \leq \frac{2(s-1)}{s^2}$$

whenever the optimal coefficient is formulated as an infimum of such uniform expansions. For $s = 16$ these bounds are 23/512 and 15/128. Can cancellation between two erroneous derivatives close the gap?

[write] **7 October 2026; sources 60–69.** *Both ends improved by source 60:* $(s-1)/s^2 \leq c_s^{\text{prime}} \leq (3s-2)/(2s^2)$, at $s = 16$ from [23/512, 15/128] to [15/256, 23/256] (its Theorems II.336.2 and II.336.3); the value is open (its Research question V.226.6).

Research question V.127.2 (The next coefficient for higher arrangements). For $s \geq 6$, determine the optimal cubic coefficient after the two sharp all-group coefficients have been fixed. The full bound F_s , the simpler quadratic inverse, and the binary extremal family give different third-order information. A useful first target is $s = 6$, where one can classify low-weight error patterns while retaining the missing sixth independence degree exactly.

[write] **7 October 2026; sources 60–69.** *Answered by source 60:* the cubic coefficient is $8(s-1)(2s-1)/(3s^3)$ for $s \geq 6$, and every further coefficient is optimal in succession (its Corollary II.335.3).

Research question V.127.3 (Equality and near-equality in the sharp cubic). Classify the maps with $0 < \delta < 1/4$ and $\varepsilon_4 = 4\delta - 24\delta^2 + 32\delta^3$. The proof forces simultaneous equality in the modal collision estimate, the increment cap, the missing-moment bound, and the one-error implication. Do these conditions characterize binary row-defect constructions, perhaps after quotients, products, and addition of an exact model? A quantitative near-equality theorem would identify the actual obstruction to a better repair modulus.

[write] **7 October 2026; sources 60–69.** *The equality part is answered by source 60* for every even $s \geq 4$: equality at $0 < \delta < 1/4$ holds exactly for $a\mathbf{1}_C(x)\lambda(y)$ plus an exact model, C a coset of index $n \geq 3$, at $\delta = 1/(2n)$ (its Theorem II.334.3); near-equality is its Research question V.226.5.

Research question V.127.4 (The activation threshold and the global distance spectrum). How far can the small-branch theorem extend? The bound $1/(240(s-1))$ was chosen to make the elementary bootstrap automatic. A sharper noisy-affine lemma, or a direct spectral gap argument, may increase this threshold substantially. Separately determine the feasible pairs (δ, ε_4) away from zero. The vanishing of F_4 at $1/4$ shows why extrapolating its inverse without a structural argument is not a solution to the global problem.

Research question V.127.5 (A sharp theorem from row profiles). Theorem II.140.3 retains $\bar{t}/\delta$, while the uniform-row cap retains only $\max_x \delta_x$. What is the optimal defect bound given the complete row distributions p_x , or just the first few moments of δ_x ? Such a result could distinguish diffuse errors from concentrated row defects before passing to a worst-case constant.

Research question V.127.6 (Sparse domains with quantitative mixing). Replace near-full density by explicit assumptions that conditioned arrangement queries have approximately uniform marginals and that increment cocycles have enough common witnesses. Can one prove $\delta_A \leq C\eta$ with polynomial dependence on these mixing parameters, instead of paying the invalid-arrangement error $2s\tau$? The precise challenge is to replace the three full-product steps: auxiliary-label independence, mode comparison at every increment, and rowwise integration. This is the most direct route toward the sparse setting of Gowers’s Sections 12–13.

Research question V.127.7 (Nonzero increments and nonrepeated vertices). Determine exact conversion bounds when the test excludes $h = 0$ or forbids repeated vertices. Conditioning changes both the total count and the independence law. The crude invalid-event argument gives an initial estimate, but a sharp normalization should distinguish genuine collisions from repetitions forced by the balanced relation. In small characteristic, these distinctions may affect extremizers.

Research question V.127.8 (Two-direction testing and the row gauge). Add the horizontal analogue of the arrangement test. The vertical test alone cannot constrain $a(x)$, while a horizontal test should force compatibility with an affine function of x . Establish a quantitative repair theorem to a fully biaffine map, with explicit dependence on the two failure rates and sharp first-order constants. The intersection of the two exact model classes is easy to describe; robust intersection with optimal constants is the substantive issue.

[write] **7 October 2026; sources 60–69.** *Answered locally by source 60*: the optimal intersection modulus $J(\eta)$ for $\eta < 1/4$ and biaffine repair with optimal inverse coefficients through degree $s-2$, sharp first-order constant $1/(2s)$ (its Theorems II.337.2–II.337.7); the full finite-error profile is its Research question V.226.8.

Research question V.127.9 (Higher multilinear arrangement models). For products of three or more groups, formulate the gauge class left invisible by a prescribed family of directional difference

tests. Can the coarse derivative synchronization be iterated without an exponential accumulation of error constants? The multilinear structure problems in [129] and directional norms in [130] provide relevant settings, but the exact high-agreement tester and its nullspace should be fixed before estimating constants.

Research question V.127.10 (Fast decoding rather than exhaustive slopes). The exact cyclic decoder avoids exponential enumeration of row functions but still enumerates all permitted (b, c) . Can random derivative queries identify these coefficients in near-linear time in the table size, under the high-agreement promise? A robust decoder should output both a model and an efficiently verifiable disagreement count; it should also expose its sample complexity and failure probability rather than hide them in an informal majority vote.

Research question V.127.11 (A formally checked probabilistic core). Formalize the auxiliary-label correction lemma, affine separation, proper-coordinate independence, and the one-error inclusion–exclusion identity. These components are small enough to isolate from the full Szemerédi development. Once checked, they would certify the main theorem without requiring a formalization of a general inverse theorem for Gowers norms. The initial target should be the rational inequality $\varepsilon_s \geq s\delta - 2s(s-1)\delta^2$ on the proved small branch, not a real square-root expression.

Research question V.127.12 (Actual propagation through the Gowers pipeline). Find a point in the sparse structural argument where a near-full product or a quantitatively mixing product can be produced without losing the information needed later. Only then compare the complete parameter recurrence with the existing one. A successful answer must keep track of domain density, valid-arrangement mass, failure rate, model agreement, and the eventual progression length. This is the necessary bridge between the present local refinement and a claim about a global Ramsey or Szemerédi bound.

The most immediate mathematical targets are the first and third questions: the common-prime quadratic constant and the equality classification. The most consequential integration target is the sixth, because the present proof deliberately does not solve the sparse-domain mixing problem.

V.128 Repository integration and formalization interface

Source 33, delivered Section 10.

V.128.1 Recommended placement and source pin

A distinct research subdirectory is preferable to silently replacing an existing lemma:

```
Combinatorics/Ramsey/Research/GowersSzemerédi/  
respected-arrangement-stability/
```

The mathematical entry point is Theorem F.38.2, with Corollary F.38.4 as the eight-arrangement specialization. The article and its code do not alter the repository’s formalization ledger.

The source crosswalk was checked at repository commit

9f83dbe0f15205c10b1e6efc6af69b09969f8e03

The corrected Gowers transcription has blob identifier

41085a0efde8c1932e86e80c791984f52841c22e.

The local-refinements README at that pin has blob identifier

b0edfc23a3624a09c7a2087a336d447a4378622c.

The repository was changing during preparation; earlier exploratory reads were from the moving branch. The pinned reads above, not an assumption that every exploratory read had the same commit, support the final source crosswalk. The accompanying source audit records the specific ranges and the scope of the comparison.

[write] **6 October 2026; sources 28–33.** *Checked at HEAD 5d4718772.* The transcription blob 41085a0 is unchanged; the README of this report has changed (writes of sources 11–27). `lemma_12_3`, `lemma_12_5` and `lemma_12_6` are at lines 151, 170 and 183 of `Sections12_13.lean`; none has a counterpart for the repair theorem, and nothing here is formalized.

V.128.2 Suggested formal layers

The proof can be implemented over finite additive commutative groups, using finite sums and rational probabilities initially. There is no need to import an unproved sparse inverse theorem.

Layer	Mathematical content	Main dependency
Arrangement domain	Balanced tuples, occurrence counts, proper-coordinate independence	Finite group bijections
Labeled correction	Difference modes, exact additivity, $e/(1 - 2e)$ repair	Finite averaging and union bounds
Derivative synchronization	Affine-map group, separation, increment cocycle	Labeled correction
Local defect	Row modes, collision profile, one-error probability	Independence and finite inclusion–exclusion
Repair theorem	Small branch and quadratic inequality	Coarse and local estimates
Higher orders	Exact moment formula and monotonicity	Finite Fourier analysis
Sharpness	Cyclic models, exact example counts, series coefficients	Elementary finite algebra

Table V.128.1: A proposed formalization dependency order; no row is represented as already implemented or checked in Lean.

The cycle deficit lemma can be formalized independently as a finite probability-vector inequality, followed by its subprobability extension. The accompanying moment interpolation should retain ambient count normalization in the partial-domain statement.

For $s = 4$, all ingredients of the repair theorem can be stated without Fourier analysis. For higher even s , one can formalize the moment identity separately and reduce to that coarse decoder. The local polynomial inequality is itself Fourier-free for every s .

A formal implementation should distinguish an exact count from its normalization, a function-valued equality of affine maps from pointwise agreement on most x , and an auxiliary label sampled at each occurrence from a fixed random section. The last distinction is especially important for repeated coordinates. The proof also needs the positivity hypotheses $1 - 2e > 0$ and $1 - 12d > 0$ at their precise uses, rather than after clearing denominators without a sign check.

V.128.3 Claim ledger

All mathematical results proved in this manuscript are unconditional under their displayed hypotheses. The conditional statements are the erasure corollary's explicit coverage assumptions and the clearly posed future research questions. No novelty claim is used as a proof input. The majority-correction method is classical in origin, while the arrangement-specific estimates and constructions are developed here. The exact cubic and sharp asymptotic coefficients are offered as research contributions for scrutiny, not as refereed literature facts.

No new statement in this package is asserted to be Lean-compiled or kernel-verified. The supplied Python checks are ordinary exact-arithmetic programs. No new global estimate for $r_k(N)$ or van der Waerden numbers is claimed. These limitations are part of the mathematical scope, not missing hypotheses to be filled by optimism.

V.129 Normalization and audit checklist

Source 33, delivered Appendix A.

V.129.1 A compact normalization table

Quantity	Convention
Gowers arrangement order	$m = s/2$
Vertical difference occurrences	s
Point occurrences	$2s$
Free horizontal coordinates	$s - 1$
Independent vertical base points	s
Common increments	one element of H
Total full-domain arrangements	$ G ^{s-1} H ^{s+1}$
Distance	uniform pointwise disagreement on $G \times H$
Character average over target	$ K ^{-1} \sum_{\chi \in \widehat{K}}$
Frequency sum over horizontal group	$\sum_{\xi \in \widehat{G}}$

V.129.2 Degenerate inputs and boundary cases

The all-zero increment is included, even though it always passes. This matters for both the binary and the prime extremizers. Repeated vertices and repeated horizontal coordinates are also included. The auxiliary-label proof remains valid for those repetitions because labels are resampled per occurrence. In a partial-domain test, validity means that every occurrence has a value; repeated occurrences do not require separate membership choices.

For $\delta = 0$, ratio expressions involving $\bar{t}/\delta$ are not used; exact classification supplies the zero-error conclusion. For empty valid-arrangement sets, conditional failure is undefined. For trivial target groups, every function is an exact model and every theorem reduces to zero error. The order parameter is always even and at least four; none of the signed inclusion–exclusion claims has been extended to odd orders without proof.

V.129.3 Numerical evaluation of the repair formulas

For small u , cancellation can make the expression $1 - \sqrt{1 - 8(s-1)u/s}$ numerically inconvenient. An equivalent stable form is

$$R_s(u) = \frac{2u/s}{1 + \sqrt{1 - 8(s-1)u/s}}. \quad (\text{V.129.1})$$

The exact verification suite avoids square roots entirely: it verifies that δ is on the increasing branch and that $Q_s(\delta) \leq u$. For the cubic inverse one may use bisection on $[0, 1/12]$, since the cubic is strictly increasing there. A floating-point root returned without checking the interval is not a substitute for that branch condition.

V.129.4 Reproducibility commands

From the package directory, the following commands need only a standard Python installation and a conventional pdfLaTeX distribution:

```
python code/verify.py --extended
python code/verify_erasures.py
python code/arrangements.py data/example.json --orders 4 16
pdflatex -interaction=nonstopmode -halt-on-error article.tex
pdflatex -interaction=nonstopmode -halt-on-error article.tex
pdflatex -interaction=nonstopmode -halt-on-error article.tex
```

The third LaTeX pass stabilizes the table of contents if page numbers changed on the preceding pass. The package README and Makefile supply the same routes. All mathematical labels in this source use the prefix **ras:** to reduce collisions when integrating with a larger article; macros and theorem environments still need the normal editorial merge.

V.130 Consequences, limitations, and a research program

Source 34, delivered Section 11.

The results above settle several questions in the audited repository: the exact values $h(7), h(11)$, the leading asymptotic of $h(p)$, the support cost and classification of positive cyclic detectors, and the proposed single-pair model for the fourth-norm extremizer. Appendix II.163 settles the strict threshold in the nonuniform directional model. The following questions identify the remaining mathematical gaps. They are proposed research targets; we do not assert that every formulation is historically new or unresolved throughout the literature.

V.130.1 Binary determinants and arithmetic configurations

Research question V.130.1 (Divisibility versus exact cyclic isolation). Is $h(p) = h_{\text{cyc}}(p)$ for every prime p ? If not, what is the first prime where they differ, and how large can the difference be?

Theorem III.75.2 concerns a determinant divisible by p , while $h_{\text{cyc}}(p)$ requires that the complete row quotient be exactly $\mathbb{Z}/p\mathbb{Z}$. These conditions agree in every small case proved here and in their leading asymptotic, but neither conclusion proves exact equality in general. A useful finite search should store full row-lattice indices and full-support characters, not only one determinant. Existing determinant spectra can supply many immediate values before any new exhaustive search is attempted.

Research question V.130.2 (The second-order support term). With $L = \log p$, $\lambda = \log L$, and $\mu = \log \lambda$, does there exist a constant a such that

$$h(p) = \frac{2L}{\lambda} \left(1 + a \frac{\mu}{\lambda} + o\left(\frac{\mu}{\lambda}\right) \right)?$$

Ask the same question for $h_{\text{cyc}}(m)$ as m runs through all integers. Can the coefficient 3 in the proved upper comparison be reduced toward the coefficient 1 in the Hadamard comparison?

The bounds here place any such second-order coefficient between one and three; they do not establish its existence. Improving the consecutive determinant range and improving the representation of one specified prime are different problems. The latter could admit a better upper bound even if an all-integer interval is more costly.

Research question V.130.3 (The secondary loss in the consecutive cyclic range). Can every positive integer up to

$$\exp\left(\frac{1}{2}N \log N - o(N \log \log N)\right)$$

be realized as the sole nonunit Smith factor of a binary matrix of order N ? More modestly, can the coefficient of $N \log \log N$ be made strictly smaller than one?

In the present construction, the bounded first row is encoded at a cost proportional to $\log A$. Balancing that cost against the number of radix positions leads to an auxiliary order of scale $N/(\log N)^2$. This explains the current secondary loss within this construction. A smaller loss would likely need a more efficient use of the first row or a different composition mechanism. The cyclic condition should be retained throughout; a determinant construction alone does not resolve the stronger question.

Research question V.130.4 (Exact positive support and finite-size efficiency). Determine $h_{\text{pos}}(m)$ exactly for a substantial range of cyclic orders. When is the doubling bound $h_{\text{pos}}(m) \leq 2h_{\text{cyc}}(m)$ strict, and what is the second-order term in $h_{\text{pos}}(m)$?

The leading constant four and the inverse-pair classification are proved here. The remaining finite problem is closely related to ternary matrices with a prescribed cyclic quotient. The doubling construction gives a simple bridge from binary matrices, but it need not minimize the number of pairs. A useful search should retain a full-support faithful quotient vector and certify the whole lattice index. The corresponding positivity classification for noncyclic finite quotients is a further structural question.

Research question V.130.5 (Prescribed finite abelian quotients). For a finite abelian group A , determine the least size of an affinely independent Boolean configuration with row quotient A and with every coordinate generator nonzero. How does this size depend on $|A|$, the number of invariant factors, and their prime power structure?

The cyclic construction controls a single nonunit invariant factor. Direct sums of binary presentations give an upper bound for general A , followed by affine augmentation and deletion of coordinate generators that vanish in the quotient. The substantive question is whether several invariant factors can share the same Boolean coordinates efficiently. Prime powers also distinguish a rank drop modulo p from its full p -adic depth.

Research question V.130.6 (From first occurrence to a complete coefficient). Classify all full-support lattice types at the first few unresolved support degrees, and compute their multiplicities as functions of the cube dimension. In particular, can the sixth- and seventh-order coefficients be expressed in a finite, explicit basis of Fourier statistics with a dimension-uniform remainder?

The onset theorem supplies necessary dimension and support thresholds, and the explicit matrices supply terms that must be present. They do not enumerate all rational dependencies, all Smith types, or all cancellations among supports of the same size. A practical approach is to classify coordinate-pattern orbits, attach integral presentations to each orbit, and then derive closed multiplicity formulas. A useful theorem must keep the distinction between Smith torsion and reductions of pre-existing integer relations.

V.130.2 The actual Gowers norm extremizers

Research question V.130.7 (The sharp odd fourth-norm constant). Determine c_4^{odd} inside the proved interval

$$\left[\frac{4804}{5465198316262956^{1/4}}, \frac{2}{3} \right].$$

Is the supremum attained on some finite group, or only approached along groups of growing order? What are the extremizing Fourier supports and phases?

The dual identity identifies why the single-pair guess fails, and the resonance argument controls a concrete way to exploit that failure. It does not classify genuine minimizers. On a fixed group, the finite-dimensional variational equation in Remark III.71.8 is a necessary condition. Across all groups, an additional compactness or approximation argument would be needed. Merely optimizing the two-pair family does not establish a global optimum.

Research question V.130.8 (Stronger coupled certificates). Can the dominant pair and its resonant harmonics be retained jointly through all positive autocorrelation terms and the dual inequality, so as to improve $2/3$? Can a finite family of explicit dual tests provide a sharp certificate?

The proof uses only the balanced monomials and the families with seven first harmonics and one third or fifth harmonic. Further families impose additional positive costs on the same Fourier mass. Their phase-sensitive dual contributions may also be useful. A sound computational program should first produce exact polynomials and rational inequalities, then prove the resulting optimization bound. Floating-point minimization by itself supplies a candidate, not a universal norm inequality.

The displayed two-pair polynomial gives a particularly accessible subproblem. Its nonzero stationary parameters satisfy

$$(1 + t^4)P'(t) - 16t^3P(t) = 0.$$

Exact root isolation can determine the best member of that family. One must then test perturbations in additional harmonics to decide whether a family optimizer is even locally extremal in the full function space.

Research question V.130.9 (Higher dimensions and quantitative stability). Determine the asymptotic behavior of c_d^{odd} as $d \rightarrow \infty$, and develop stability statements for near-minimizers of $\|f\|_{U^d}/\|f\|_{U^2}$ for $d \geq 4$.

The sequence is nonincreasing by norm monotonicity. The incoming moment estimate and the profile refinement in Remark III.71.10 give explicit upper bounds, but the fourth dual function stops improving the moment endpoint in sufficiently large dimensions. Computing $\mathcal{D}_d$ on a small Fourier support becomes a Boolean balancing problem related to the configuration analysis in the first half of the article. Stability must be formulated for the genuine norm minimizers; stability of an intermediate autocorrelation inequality has a different equality set and is not enough.

V.130.3 Returning to the density-increment argument

Research question V.130.10 (A constrained density-and-norm envelope). For fixed $0 < \delta < 1$ and small u , determine or sharply bound

$$\sup\{Q_4(F) : 0 \leq F \leq 1, \mathbb{E}F = \delta, \|F - \delta\|_{U^4} = u\}.$$

Can the quartic contribution and the higher-order remainder be bounded jointly using their compatible Fourier phases?

The coefficient $100c_4^{\text{odd}}$ describes the unrestricted small-amplitude leading term for real perturbations. The box constraint $0 \leq F \leq 1$ becomes more restrictive at larger amplitudes. In addition, the individual support bounds used in the remainder need not be simultaneously sharp. A joint envelope could therefore improve the useful finite-error range even before the exact value of c_4^{odd} is known.

Research question V.130.11 (Propagation through a complete iteration). Recompute a specified, fully normalized density-increment scheme after replacing the local fourth-dimensional cube bound and the directional rigidity threshold by the results proved here. Which final constants or exponents actually change?

The appropriate calculation must track density, progression or ambient-set size, uniformity thresholds, and all recursive dependencies. A smaller local coefficient may affect only a final constant; in some parameter regimes another loss may dominate it entirely. The current article supplies explicit replacement inequalities but makes no unverified claim about their global effect. This is also the setting in which the prime-order and nested hypotheses of the original inverse argument must be preserved.

V.131 Reproducibility and proposed ProveIt integration

Source 34, delivered Section 12.

V.131.1 What is proved mathematically and what is checked finitely

The statements for every prime, every cyclic order, and every finite odd-order group are proved by algebraic and analytic arguments in the article. Their validity is not extrapolated from a numerical search. The finite computations have three roles: certifying specific matrices, supplying an exhaustive negative certificate for two small primes, and verifying exact coefficients or rational comparisons in the analytic proof.

Component	Exact verification	Role in the proof
Small prime supports	Determinants and full-support syndromes for 7, 11, 13, 17	Positive finite certificates
Seven-vertex exclusion	All 104 and 204 projective orbit representatives; every Boolean row and integer annihilator	Exhaustive negative proof for 11, 13
Positive cyclic construction	36 moduli; two direct spatial moment checks	Implementation check of an all-parameter proof
Determinant construction	144 bounded-row encodings and 92 complete compositions; exact determinants, Smith forms, and unit cofactors	Check of signs, dimensions, and rounding
Fourth norm	Dual counts, resonant coefficients, order-three polynomial, rational $2/3$ inequalities	Exact confirmation of identities proved in the text
Two-pair example	All 4^8 face assignments and exact evaluation at $-1/7$	Finite coefficient certificate

Table V.131.1: The verification methods and their logical roles are distinct.

The code and data are included in the archive. The torsion and norm verifiers use Python’s standard library. The determinant gadget verifier uses exact SymPy arithmetic; the dependency is recorded in `requirements.txt`. No floating-point result is needed for a theorem stated here. Decimal values in the text are provided for orientation; the defining expressions and decisive inequalities are exact.

The negative certificates deserve a particularly clear trust boundary. Their generator uses rational elimination to find an annihilator, but the verifier checks the supplied nonzero integer vector directly against every admissible Boolean row. The verification does not trust a rank routine. Exhaustiveness comes from an explicit enumeration of sorted tuples and all nonzero scalar multiples. This structure is suitable for a later formal certificate checker.

V.131.2 Source control and placement

The suggested new location is a research directory beside the existing merged report:

`Combinatorics/Ramsey/Research/GowersSzemerédi/prime-torsion-norm-gaps/`.

The package contains a single self-contained TeX source, the compiled PDF, a build file, exact verification scripts, certificate data, and a source ledger. It can be placed independently without renumbering statements in the existing merged article. A later editorial merge should retain source attribution and stable theorem labels.

The inspected formal baseline is `Proofs03CubeUpper.lean`, in the existing `Combinatorics/Ramsey/Lean/GowersSzemerédi/` directory. At the pinned commit, the status ledger records `lemma_3_10`

as `exact_companion`, with companion `lemma_3_10_holds`. That existing proof was inspected. This report does not supply new Lean proofs or report a fresh repository build.

[write] **6 October 2026; sources 34–39.** *Checked at HEAD 89c1cbc38.* `lemma_3_10` is at `Sections01_03.lean`, line 340, with the exact companion `lemma_3_10_holds` (`Proofs03CubeUpper.lean`, line 82; the file has the blob `c017515` the source records). No declaration concerns c_d^{odd} , prime torsion of cube expansions or the transport model of Appendix A; nothing of this source is formalized. The suggested separate directory was declined at placement (`bf2fe146d`): the source answers questions already open in this report.

V.131.3 A concrete formalization sequence

1. Define the augmented Boolean matrix and its integer row quotient. Formalize the Fourier constraint identity and the link between rank drop modulo p and Smith torsion.
2. Prove the affine-minor reduction and the minimal-order full-support lemma. These require finite-dimensional linear algebra and elementary determinant identities, without the Alon–Vu theorem.
3. Add the small explicit matrices and the finite certificate checker. The $p = 11, 13$ lower bounds can then be checked using integer dot products and finite orbit enumeration.
4. Formalize the positive cyclic presentation and cyclic support compression. State the Alon–Vu input separately before formalizing the determinant gadgets and asymptotic corollaries. No unproved external input should be hidden behind a theorem named as if it had been established internally.
5. Formalize the exact dual polynomial, positive balanced and resonant moment sums, and the order-three calculation. Derive the rational $2/3$ bound from these ingredients, then connect it to the existing normalized cube-count interface.
6. Formalize the finite-fiber sampling model in Appendix II.163. Keep its general-group sharpness theorem separate from any stronger claim at the original prime-order inverse-theorem interface.

The most valuable initial formal additions are the exact reduction, the small prime certificates, and the fourth-norm inequality. They have compact inputs and conclusions, and do not require formalizing the entire asymptotic matrix construction first.

V.132 Evidence, source scope, and formalization

Source 35, delivered Section 8.

Delivered part heading of source 35: **Verification, integration, and further research**

V.132.1 What the accompanying calculations check

The mathematical proofs above establish the infinite families of statements. The programs serve a narrower purpose: they reproduce finite identities, verify the exact numerical certificates, and test the delicate conventions on complete small examples. They use the Python standard library only. No floating-point optimization is an input to a theorem.

Arrangement selection. The program `verify_graded_selection.py` examines all $3^8 = 6,561$ coefficient arrays for $k = 1, m = 4, p = 3, L = 2$. It removes the universal line, computes the parity-removed degree, and solves the resulting linear feature systems exactly over $\mathbb{F}_3$. The maximum success probabilities are

$$Q_0(3) = \frac{1}{3}, \quad Q_1(3) = \frac{11}{81}.$$

These maxima attain the stated individual-certificate estimates. The program also checks the dimensions of balance subspaces on separate small examples and records exact rational product masses. An attained certificate probability does not imply sharpness of the full selection threshold.

For a direct check of actual subset selection, the same program takes $B = \mathbb{F}_3^2$ and $\phi(x, z) = x^2z$, enumerates all 6,561 arrangements, and averages over all 243 field seeds. It uses products over distinct physical points for the actual probabilities and products over labels for the Fourier comparison. In this example there are 3,753 respected and 2,808 unrespected arrangements. The difference between the expected actual and labelled bad counts is

$$\frac{9,891}{128} - \frac{84,687}{2,048} = \frac{73,569}{2,048} > 0.$$

Thus omitting the repeated-vertex compensation would be a real error, not merely a missing discussion of an empty exceptional set. The program further verifies 54 exact rational instances of the explicit bandwidth budget and recomputes the exponent table through $k = 16$.

The fourth-order norm gap. The program `verify_norm_gap.py` checks all 17 rational comparisons used in the scalar certificate. Comparisons involving fourth roots are reduced to powers of positive rational numbers. It also enumerates all 2^{15} sign assignments for the cosine dual and all 2^{16} assignments for the full order-three cube. The resulting coefficient dictionaries are

$$\{-5 : 1, -3 : 27, -1 : 111, 1 : 111, 3 : 27, 5 : 1\}$$

and

$$\{-12 : 8, -6 : 16, 0 : 286, 6 : 16, 12 : 8\},$$

respectively. These checks support the combinatorial proofs of Lemmas III.72.3 and III.72.4; the analytic energy and scalar arguments remain explicitly proved in the text.

Polynomial obstructions. The program `verify_polynomial_obstruction.py` tests 285,915 starting-residue and step-residue pairs in the three examples tabulated in Part III. Arc diameters are computed by the largest-gap formula on integer residues, so endpoints and wraparound do not involve rounding error. It also tests all 422,500 length-three integer progressions in the 1,301-point interval of the prime-modulus approximation, using the prime

$$N = 54,121,601.$$

At tolerance $1/32$, the good steps in that interval are exactly 217 and 434, both divisible by the predicted period 217. The prime test is local to this interval.

Reproduction. The archive includes recorded JSON outputs and a runner, `verification/verify_all.py`, that reruns the suites into a temporary directory and compares their outputs with the records. It leaves the supplied evidence files unchanged. The $\text{\LaTeX}$ source contains its own bibliography and does not depend on external graphics, a bibliography database, or the original repository being available at build time.

V.132.2 Which ingredients are inherited

The arrangement architecture, the respected relation, and the use of positive Fourier expansions originate in Gowers’ argument [1, 158]. The one-stage Fejér selector and its polynomial retained-count powers already appear in [160]. The degree-sensitive certificate geometry is developed in [161]; we reprove it for the arbitrary field coefficients needed here. The contribution is to place the necessary top-degree balances *inside the probability-mass estimate*, combine degree and codimension before summing, and carry the physical-collision correction through that weighted calculation.

The real centered U^3 comparison and the centered cube expansion are in the merged collection [159]. The fourth-order cosine identities and the first uniform higher-order norm gap are in [162]. The strongest inspected fourth-order comparison before this article is the moment coefficient $35/8$ in [163]. Our additional positive Fourier configurations increase that coefficient to 5. At higher orders the existing moment theorem remains stronger than the constant obtained just by propagating 5.

For polynomial partitions, the linear divisibility-and-conductor argument is in [165]. The preceding polynomial volume obstruction is in [164]. The present descending-degree construction supplies a divisor that is forced at every initial point; this is the additional property needed to multiply the polynomial obstruction by the semigroup conductor.

V.132.3 The literature comparison and its limits

The foundational context is Gowers’ 2001 proof [1]. The global progression-free bound cited in the introduction is from Leng, Sah, and Sawhney [2]. Neither the size threshold of a local restriction nor an improvement of its numerical exponent can be identified with an improvement of that global bound without a further complete argument.

On the analytic side, Eisner and Tao study large values and near-extremizers of Gowers and Gowers–Host–Kra norms [7]. Kovač and Rogers study L^p variants, degree-lowering estimates, and strict log-convexity [71]. Their Theorem 1.1 assumes a nonnegative function and a group without compact open subgroups; its stated strict log-convexity conclusion consequently does not apply to finite groups. Our comparison concerns signed real centered functions on finite odd-order groups and relates U^4 directly to U^2 . These are different hypotheses and inequalities. This distinction is useful context, but is not an exhaustive priority search.

For polynomial phases, Green and Tao’s quantitative work on polynomial orbits provides broader context for smoothness estimates [148]. The small-arc lemma used here has its own finite-difference proof and does not invoke a nilmanifold equidistribution theorem.

The source audit therefore supports precise comparisons with the inspected statements. It does not certify that every intermediate lemma or the final combination is absent from all earlier literature. The manuscript is intended for mathematical review and integration as a research contribution, with its proofs available for independent checking.

V.132.4 A concrete formalization route

The repository’s statements and their formal companion proofs must be distinguished. We inspected the ledger and the relevant section files at the source pin; no new Lean file is included or claimed to compile. The present theorems are conventional mathematical proofs, supported by exact computations. Table V.132.1 gives separable proof obligations for a future formal development.

[write] 6 October 2026; sources 34–39. Checked at HEAD 89c1cbc38. lemma_15_5, lemma_14_8

Component	First formal target	Main dependencies
Weighted selection	Product mass of an affine subspace; dimension D_t	Finite product probability, coordinate projections, Boolean multilinear interpolation
Certificate geometry	The t and $2t$ minors, including constant levels	Finite-field rank and counting; independent linear forms
Subset extraction	Expected labelled counts, collision correction, score outcome	Additive-character orthogonality; a distribution on finite subsets
Real norm gap	The energy inequality with $J(\lambda)$	Normalized Fourier identities; finite group parity and torsion facts; cube Cauchy–Schwarz
Scalar constant	The rational two-interval certificate	Positive fourth roots, differentiation, rational arithmetic
Polynomial obstruction	Short-arc lifts and descending divisibility	Finite differences, interpolation, prime divisibility, numerical semigroups

Table V.132.1: Formalization tasks; no row is presented as completed Lean work.

and `lemma_5_9` (`Section05.lean`, line 261; companion `lemma_5_9_holds`, `Proofs05FullPartition.lean`, line 98) are exact; `lemma_3_10` is exact (companion `lemma_3_10_holds`, `Proofs03CubeUpper.lean`, line 82). No statement of this source is formalized; the transcription it cites has the same blob 41085a0 at HEAD.

For integration with the existing normalization, the first interface checks should be the ordered-arrangement count q^M , the density upper bound (II.145.2), and the conversion between normalized and unnormalized Fourier coefficients. For the partition result, encode a proper progression with an integer lift and explicitly quantify over all target lengths as in Corollary I.61.6. Those choices avoid apparent numerical discrepancies caused solely by differing conventions.

V.133 A proposed research programme

Source 35, delivered Section 9.

The following questions are ordered by their proximity to the proved arguments. Each has a concrete first target. Questions about optimality are proposed problems, not consequences of the bounds already established.

V.133.1 Arrangement selection: beyond the present exponent

Research question V.133.1 (Top-degree rank, not only top-degree size). Can the coefficient strata be refined by the rank of the matrix

$$(A_{j,T})_{\substack{1 \leq j \leq m \\ |T|=t}}$$

so that several lowest Boolean moments contribute independent restrictions simultaneously?

The present proof selects one nonzero moment vector. Multiple independent vectors could force more restrictions on both the levels and the anchors, while their rank constraints also reduce

coefficient mass. A meaningful first result would compute the joint rank-and-mass profile for $k = 2$ or 3 , prove its exceptional varieties have the predicted dimension, and improve the explicit exponent 23 for $k = 2, m = 16$. Summing a list of pointwise rank bounds without controlling the mass of each rank class would not suffice.

Research question V.133.2 (Exact mass of the arrangement balance spaces). How close is the product Fejér law to the upper bound μ^{s-D_t} on $\mathcal{V}_t$? Can its exact additive convolution structure replace the largest-atom estimate?

The mass lemma is sharp for arbitrary coordinate subspaces, whereas $\mathcal{V}_t$ comes from many Boolean evaluations and balanced sums. Even an exponential-in- s constant improvement would affect practical thresholds. A changed power of L would require proving that the particular support intersects these spaces less efficiently than a general subspace. An anisotropic version could assign different bandwidths to cube positions, provided the respected universal signal remains large enough.

Research question V.133.3 (Optimize the kernel against the whole error budget). Among nonnegative trigonometric polynomials $K \leq 1$ with nonnegative Fourier coefficients, which best balances universal signal, coefficient-stratum mass, and point marginal probability?

The earlier positive-Fourier report studies the signal and its retained-mass powers [160]. The present theorem adds a different objective: the entire graded error budget. A kernel that maximizes the scalar ratio $\sum c_n^s/c_0^s$ need not minimize $\sum_t q^{-t-1} \Pr(b \in \mathcal{V}_t)/c_0^s$. A finite-bandwidth optimization for $k = 2, m = 4$, followed by an exact certificate for the optimizer, is a manageable first problem.

Research question V.133.4 (Small characteristic and torsion in the target). Can one obtain a comparable theorem when the characteristic is fixed while $\alpha\eta$ tends to zero, or when the target has small torsion?

The current restriction $p > 2(L - 1)$ serves a definite purpose: the integer Fourier band must embed without collisions, and nonzero universal frequencies must not annihilate a defect. Simply replacing q by a large extension-field cardinality does not remove this restriction. A possible approach is a multidimensional kernel on the dual of a larger auxiliary group, with enough linear maps to detect every nonzero target defect. The corresponding universal signal and mass estimates would have to be proved anew.

Research question V.133.5 (Intrinsic lower bounds for selection). What obstruction forces a lower bound on the ambient field size for any subset satisfying both the desired purity and a retained count of order $\alpha^s \eta^{s-1} T_0$?

The asymptotic $E_{k,m} \sim 2s/k$ describes our sufficient threshold, not a matching necessary one. A serious lower bound would need a family of maps and domains defeating *all* candidate subsets, not just one kernel or one relation-counting proof. It would be useful first to separate optimality within positive-Fourier selectors from optimality for unrestricted subset selection.

V.133.2 Norm gaps: sharpen the analytic certificate

Research question V.133.6 (The exact fourth-order extremal constant). Determine c_4^{odd} , or narrow substantially the interval in (III.72.7). Is a maximizing sequence supported on a bounded number of Fourier frequencies?

The retained statistic $J(\lambda)$ sees every amplitude three-term progression, while the proof of the coefficient 5 uses only a small part of it. A first improvement could retain all terms involving the dominant pair and its third and fifth harmonics, then prove an exact polynomial inequality on their amplitudes and phases. Numerical optimization can suggest a candidate but must be followed by a uniform certificate and a treatment of character-order coincidences.

Research question V.133.7 (Quantitative structure near the minimum). If $\|f\|_{U^4}^{16} / \|f\|_{U^2}^{16}$ is near its infimum, what Fourier concentration and phase alignment are forced?

The inequalities $R \geq (5 - 3p)^2$ and (III.72.17) already restrict the largest pair and its third harmonic. They do not yet classify near minimizers. A useful stability theorem should identify a metric, quantify its distance to a candidate family, and be uniform over all finite odd-order groups. It must also accommodate the non-cosine examples already present in [162].

Research question V.133.8 (Higher-order gains beyond central-binomial moments). Can the full cosine dual D_{dg} , together with higher positive Fourier configurations, improve the established central-binomial comparison for $d \geq 5$?

The target here is a new coefficient that actually beats [163] at the specified order. Monotonicity from our U^4 result is insufficient. A tractable first step is to derive and exactly verify the dominant harmonic terms of D_{5g} , identify which Fourier energy terms penalize their negative directions, and separate the finitely many exceptional small character orders.

V.133.3 Polynomial partitions: necessity versus sufficiency

Research question V.133.9 (A matching quadratic upper theorem). For ν arbitrary quadratic phases, is the universal prescribed-length threshold bounded by

$$C_\nu L^{3\nu+2} \varepsilon^{-2\nu}$$

up to logarithmic or subpower factors, in an appropriate integer-interval formulation?

Our construction proves the corresponding lower power scale. A matching upper result would have to combine simultaneous quadratic flatness with exact coverage by lengths $L, L + 1$. Finding one long flat progression does not solve the coverage problem. A useful intermediate goal is an upper theorem allowing a small exceptional set whose dependence on L, ε, ν can be compared with the exact gap formula (I.61.10).

Research question V.133.10 (Prescribed length width and exceptional mass). How does the best universal threshold change when the allowed lengths are $L, \dots, L + w$, and how much does allowing an exceptional set reduce it?

The constructed phases give the exact conductor $L[(L - 1)/w]$ and exact uncovered mass. This provides a two-parameter family of lower benchmarks. The first upper target should interpolate between fixed w and w comparable with L , rather than treating a prescribed pair and an arbitrary-length partition as unrelated problems.

Research question V.133.11 (Prime moduli of comparable size). Do the same polynomial obstruction powers hold when the ambient modulus is prime and is constrained to be comparable with the interval length?

The periodic example is exact at a composite modulus $N = R + 1$; the prime approximation permits a much larger N . A construction respecting both primality and $N \asymp R$ would need to replace exact divisibility by a robust local condition. Conversely, a better upper theorem in this regime would show that the modulus constraint has genuine content.

V.133.4 Integration and constructive extraction

Research question V.133.12 (Quantitative propagation through the complete proof). What bound follows when the graded selection theorem is inserted into the entire corrected Gowers argument and every subsequent parameter choice is optimized consistently?

Corollary II.148.3 provides one completed substitution. The next stage should be a dependency table carrying the domain density, respect fraction, field size, characteristic, model dimension, and purity through every use. A claimed global improvement should be stated only after that propagation is complete, and then compared directly with the best applicable published theorem.

Research question V.133.13 (An effective subset algorithm). Can the expectation proof be converted into an efficient algorithm that finds B' with an explicit running-time bound in the natural input model?

The proof already yields a finite randomized construction, and exhaustive search would find a successful outcome. Efficient derandomization is a different question: computing the score $\eta X - Y$ itself can dominate the cost. A first result could use a conditional-expectation procedure for fixed k, m , with a sparse representation of ϕ and an explicit arrangement-counting oracle. Its guarantee should include both the theorem's parameters and the algorithm's complexity.

Research question V.133.14 (A reusable formal library of weighted configurations). Can the common proof components be formalized once so that future refinements change only a rank estimate, a coefficient mass, or a finite energy identity?

The product-law subspace lemma, the score-extraction argument, and the conductor packing formula are independent of most of Gowers' proof. Establishing them as reusable lemmas would make later improvements easier to check. Exact computation could then discharge the finite cosine identities and rational inequalities, while the proof assistant checks the general probability and algebraic arguments connecting them.

V.134 Computational checks and reproducibility

Source 36, delivered Section 11.

The accompanying `checks/verify.py` is an executable audit of normalizations, endpoint constructions, finite tensor cases, and boundary examples. It uses a fixed seed, 20261006. The proofs above do not rely on a numerical search or an unverified pattern.

The supplied run used Python 3.13.5 and NumPy 2.3.5. Integer tensor calculations are exact modulo the specified prime. Other tests involving complex exponentials use double precision and tolerance 3×10^{-10} ; they are not exact cyclotomic or formal verifications. The machine-readable output is `checks/results.json`.

Check	Observed result
All 81 bilinear matrices on $\mathbb{F}_3^2$; one quadratic unit-modulus extremizer for each	All attained 3^{-r} ; maximum absolute error 4.45×10^{-16}
372 mixed-function tests in dimensions 2 and 4	No violation of the alternating-rank bound within tolerance
All $3^9 = 19\,683$ functions $\mathbb{F}_3^2 \rightarrow \{1, e_3(1), e_3(2)\}$ for $B(h, x) = h_1 x_2$	Maximum $1/3$; 324 exact maximizers
All $3^8 = 6\,561$ trilinear tensors on $\mathbb{F}_3^2$, with $f = 1$	81 fully symmetric; nonsymmetric maximum $45/81 = 5/9$; 120 endpoint tensors
Both slice-profile bounds for every tensor in that census	Passed by exact integer comparison
256 tensor profiles with one seeded nonconstant bounded function	Both amplitude-sensitive bounds passed
18 432 sampled polarization tuples in characteristics 5 and 7, degrees 3, 4, 5	Exact modular identities passed
12 full $\mathbb{F}_5^2$ pointwise spectral-shift tests	Maximum pointwise error 4.86×10^{-17}
All 81 tuples in the $\mathbb{F}_3$ nonclassical example	Exact triple differences modulo 9 passed

Supplementary checks also verified both profile directions for 256 seeded multiaffine selectors with affine offsets. A direct model of $\mathbb{F}_9 = \mathbb{F}_3[t]/(t^2 + 1)$ checked 24 mixed bilinear cases, the bilinear endpoint $1/9$, and the exact trilinear endpoint count $1377/6561 = 17/81$.

The trilinear census is particularly useful for detecting the wrong threshold. A tensor has eight coefficients; its selector is evaluated on all 81 ordered pairs of increments. With $f = 1$ its energy is the number of pairs giving the zero frequency divided by 81. Symmetry is tested against swaps of the final slot with each earlier slot. Slice ranks in dimension two are zero or two and can be detected by the off-diagonal alternating coefficient. Thus the entire census, including its profile inequalities, can be checked with integer arithmetic and no Fourier approximation.

The phase-function census is also exact. If a Fourier sum has a, b, c occurrences of $1, \omega, \omega^2$, where $\omega = e_3(1)$, its squared modulus is $a^2 + b^2 + c^2 - ab - ac - bc$. The script uses this integer formula; its output includes the histogram of all energy numerators over the common denominator 729. It exhausts this finite family, not all complex-valued bounded functions. The polarization tests sample tuples and tensors, even though the computations on each sample are exact. These distinctions are recorded in the JSON and should be preserved in any integration summary.

To reproduce the checks and the document from the package root:

```
python checks/verify.py --output checks/results.json
latexmk -pdf -interaction=nonstopmode -halt-on-error sharp_symmetry.tex
```

A Makefile and a minimal dependency file are supplied. No network access is required for the mathematics or checks once NumPy and a standard LaTeX distribution are installed.

V.135 ProveIt integration and formalization plan

Source 36, delivered Section 12.

V.135.1 Placement and source fidelity

The repository comparison is pinned to 6c2e2172bab8; the full commit is

6c2e2172bab8a0af4aed6b62d53ae7cde9252e8d

The inspected corrected transcription is

Combinatorics/Ramsey/Papers/
sz-thm-gowers-proof/sz-thm-gowers-proof.tex.

Its Proposition 17.2 already has the repaired degree and characteristic hypotheses. The existing local-quantitative-refinements README describes numerous contributions on density transfer, cube expansions, floor patterns, and the local phase extraction. This manuscript focuses instead on alternating-rank certification for a global finite-field selector. This is a scoped comparison, not a claim to have exhaustively audited every pending manuscript.

A suggested independent destination is

Combinatorics/Ramsey/Research/GowersSzemerédi/
sharp-symmetry-certification/.

The label prefix `ssc:` is used throughout. Importing the text into a larger report can retain all labels, including bibliography keys. The repository has not been modified by this delivery.

V.135.2 Theorems to formalize in dependency order

The proof deliberately separates finite algebra, spectral linear algebra, and Fourier normalization. A useful dependency order is:

1. Prove Lemma IV.19.1 for finite vector-valued multilinear maps, including the empty-tuple convention. Prove the separately affine extension as a distinct statement.
2. Define the forward multiplicative derivative and normalized Fourier transform, then prove Lemma IV.16.1. Relate forward and backward derivatives by translation before connecting to an existing source predicate.
3. Formalize alternating bilinear rank, its symplectic splitting, and Lemma IV.17.2. Establish the cocycle law with its exact sign and factor $1/2$.
4. Formalize the positive operator M_g , its trace, and the invariant-eigenspace argument. This proves the main bilinear inequality without requiring the later equality classification.
5. Deduce the frozen-slice estimate and the sharp threshold. Construct the endpoint using two independent dual vectors.
6. Prove diagonal polarization and the pointwise identity. This part is independent of the spectral proof once symmetry is supplied. It is the most direct equality-form companion to the global scalar phase-extraction statement.
7. Add noise, coset, and multiaffine interfaces. Formalize the block decomposition and stability only after the main theorem is complete.

No theorem name in this plan is claimed to exist already in Lean. No uncompiled Lean skeleton is included, and no existing source statement should be marked proved on the basis of this article.

[write] **6 October 2026; sources 34–39.** *Checked at HEAD 89c1cbc38.* The transcription has the blob db4a37e the source records, and `proposition_17_2` is exact (companion `proposition_17_2_holds`, `Proofs17PhaseRemoval.lean`, line 373); `proposition_17_7` is exact (companion `proposition_17_7_holds`, `Proofs17LocalizedPhaseRemoval.lean`, line 2941). The suggested separate directory was declined at placement (bf2fe146d); the labels keep their `ssc`: stems under this report’s prefix. Formalization candidate recorded at intake, claimed by nobody: Theorem IV.17.1 for rank-two forms.

In particular, a proof of the global scalar corollary is not a proof of the local-box Proposition 17.7 or the final density increment. The companion `FORMALIZATION_PLAN.md` expands the data types and proof obligations, without inventing checked status.

V.135.3 A normalization checklist for the source interface

An integration review should check five concrete points. First, the Fourier transform in the article is normalized, so passing to the source sums introduces $|V|^{k+2}$. Second, the polynomial degree and norm order are both $k + 1$. Third, odd characteristic is used in the bilinear gauge, whereas $p > k + 1$ is only needed in classical polarization. Fourth, the input selector must be globally multilinear unless the separately affine theorem is being invoked. Fifth, the strict threshold cannot be weakened to a non-strict one. Each of these has an explicit test or counterexample in this package.

V.136 Further research questions

Source 36, delivered Section 13.

The questions below are proposed problems arising from these results. They are not all claimed to be previously unpublished, and the established low-characteristic integration results and known approximate-symmetry counterexamples remain in force.

Research question V.136.1 (All endpoint configurations). For $k \geq 2$, classify the pairs (f, σ) with T nonsymmetric and $\mathcal{A}_f(\sigma) = \vartheta_{k,p}$. Equality in the proof requires a minimum-support defect, rank two on every defective slice, full amplitude there, and a bilinear extremizer on each such slice. Which compatible global configurations occur beyond products of linear functionals and their phase/gauge variants?

The finite census in Section V.134 supplies small examples for this question, not a classification. A solution would clarify whether the endpoint is essentially a zero-selector phenomenon or whether equally extremal nonzero spectral organization exists.

Research question V.136.2 (Stability at the symmetry threshold). If T is nonsymmetric and $\mathcal{A}_f(\sigma) \geq \vartheta_{k,p} - \varepsilon$, what structural conclusion follows as $\varepsilon \rightarrow 0$? Can the defect be approximated by a product-supported rank-two defect, and can the slice extremizers be made coherent? Determine an appropriate metric before asserting a power-law bound.

Theorem IV.20.3 controls individual bilinear slices. It does not imply this global statement, because the defect support and the dependence of the frozen functions must be stabilized as well.

Research question V.136.3 (Simultaneous alternating defects). Can one improve the minimum over individual profiles in (IV.18.5) when two or more independent slot defects are present? Identify

a joint rank invariant that yields a sharp bound strictly below each single-profile estimate for a natural class of tensors.

Multiplying the separate upper bounds is not justified. Any stronger estimate must account for the fact that all slice tests use the same function and overlapping tuples.

Research question V.136.4 (Below-threshold coherent rank reduction). For energies $a \leq \vartheta_{k,p}$, can the popular low-rank slices in (IV.18.7) be organized into a structured global subspace or multilinear variety on which the remaining defect vanishes? What loss is unavoidable for a one-function derivative array, as opposed to an arbitrary multilinear form?

The intended problem retains the derivative-spectrum hypothesis. It is not the disproved unrestricted approximate-symmetry statement.

Research question V.136.5 (Arithmetic boxes and boundary losses). Develop a version of alternating-rank orthogonality for derivatives localized to arithmetic progressions or Bohr sets, with an explicit boundary error and a normalization compatible with Section 17 of Gowers’s argument. Which additional local-extension hypotheses are sufficient to turn it into a genuine refinement of Proposition 17.7?

Uniform translation of the index set was used to make M_g commute with every U_h . A local result must replace that exact step, rather than just restrict the average and reuse the global proof.

Research question V.136.6 (Lower-order multiaffine components). After Theorem IV.23.1 certifies symmetry of all nonconstant components, what is the optimal procedure for removing the lower-order frequency terms? Is there a lossless identity in an enlarged space, or an optimal quantified loss on the original space?

Example IV.23.2 rules out simply integrating the top homogeneous component and retaining the same energy.

Research question V.136.7 (Exact nonclassical certification). In small characteristic, which additional repeated-variable identities of the total form are forced by an energy threshold close to one? Can a sharp threshold certify the precise hypotheses needed for an existing nonclassical integration theorem, rather than only ordinary permutation symmetry?

[write] **7 October 2026; sources 40–46.** Answered over $\mathbb{F}_2$ in every degree by source 43 (its Theorems IV.38.2 and IV.38.4); odd small characteristic stays open (its Research question V.160.5).

The integration theorem itself is not the unresolved part of this question; the new target is a sharp spectral certification of its additional hypotheses.

Research question V.136.8 (Composite-modulus rank profiles). For finite abelian groups of odd composite exponent, express the optimal bilinear factor using the elementary divisors of the alternating bicharacter. Then seek an analogue of the multilinear support bound that yields an exact symmetry threshold and sharp examples over $\mathbb{Z}/m\mathbb{Z}$.

[write] **7 October 2026; sources 40–46.** Advanced by source 43: the bilinear factor d_B^{-1} from the elementary divisors (its equation (IV.35.9)) and the exact symmetry threshold on every finite abelian group (Theorem IV.36.4); integration on composite groups is open (its Research question V.160.6).

The finite-field support probability $1 - 1/p$ no longer describes all nonzero module maps. A useful theorem must keep this arithmetic information rather than hide it in an unspecified constant.

Research question V.136.9 (Characteristic two). Formulate and prove an equally explicit characteristic-two version of the bilinear orthogonality and spectral certification steps, including the phases required to replace division by two. Determine which parts retain the threshold $1 - (1 - 1/2)^k$ and which extra integration obstructions remain.

[write] **7 October 2026; sources 40–46.** *Answered for symmetric trilinear forms* by sources 46 and 43 (Section IV.25): exact energy $2^{-r/2}$ and codimension $r/2$; in every degree the threshold $1 - 2^{1-d}$ certifies nonclassical integrability (source 43’s Theorem IV.38.4).

The present proof assumes odd characteristic and does not claim to settle this extension. A suitable projective-operator formulation is a concrete starting point.

Research question V.136.10 (Certification from limited data). Suppose a selector is supplied by a coefficient tensor together with sampled derivative Fourier coefficients. Design a reliable test that combines multilinearity checks, a certified energy lower bound, and the robust margin $a - \eta - \vartheta_{k,p}$. Determine its sampling and arithmetic complexity without treating numerical phase evaluations as exact certificates.

Research question V.136.11 (A useful inverse-theorem regime). In which quantitative inverse arguments can the global selector be obtained with energy a and characteristic $p > k/a$ without a loss that destroys the advantage of exact certification? Track all parameters through at least one complete finite-field inverse step before claiming an improvement to a final bound.

Research question V.136.12 (A formal proof with a small analytic footprint). Can the eigenspace-multiplicity proof of the bilinear estimate be formalized using only finite-dimensional spectral theory already available in the target proof library, and then connected to the repository’s corrected global phase-extraction predicate? A useful first milestone is the exact scalar identity, followed by the vector-valued support lemma and the rank-profile theorem.

V.137 Conclusion

Source 36, delivered Section 14.

The central output is an exact extremal answer: a nonsymmetric multilinear selector on $\mathbb{F}_p^n$ can carry at most $1 - (1 - 1/p)^k$ of the normalized one-function derivative-spectrum energy, and that amount is attained. The underlying rank-profile inequality is stronger and retains information that a binary symmetry test discards. Once symmetry and the factorial-invertibility condition hold, an explicit diagonal polynomial converts the tested spectrum into a lower Gowers power with equality, even before averaging over the derivative directions.

This is a local refinement and vector-space generalization of the global phase-extraction mechanism. Its scope, counterexamples, and normalizations are explicit enough for a formalization project. Whether the sharp profile and its stability theory can be made to improve a full quantitative inverse argument is left as a separate, well-defined research task. The package supplies proofs and checks for the results stated here; it does not claim formal verification, literature priority, or an improved final Ramsey bound.

V.138 Verification and repository integration

Source 37, delivered Section 10.

V.138.1 What the accompanying checks do

The archive contains `verify.py`, its dependency list, the captured run, and `verification_results.json`. The random seed is fixed at 20261006. The run used Python 3.13.5, NumPy 2.3.5, SciPy 1.17.0, and SymPy 1.14.0. The checks have the following coverage.

Check	Number
Exact symbolic autocorrelation identities	36
Exact symbolic moments in the $R = 2$ example	2
Polynomial, matrix, root, and weight cases ($1 \leq R \leq 24$)	120
Independent discretized empty-gap minimax programs	50
Mass-minimization linear programs with contact nodes	200
Nonempty finite-cyclic indicator cases ($3 \leq N \leq 12$)	5,411
Random critical-stability measure checks	300
Exact rational recurrence-rounding checks	1,000

The largest root-modulus residual was $2.42 \cdot 10^{-14}$ after rounding up, and the largest extremal-moment residual was $8.93 \cdot 10^{-14}$. The independent gap programs used 2001 mesh points and exceeded the exact formula by at most $4.94 \cdot 10^{-7}$, as expected from discretization; the smallest difference was about $-1.90 \cdot 10^{-13}$ from solver error. The mass programs, with exact contact locations included numerically, had maximum absolute error below $1.90 \cdot 10^{-8}$.

These are regression checks, not proof certificates. In particular, floating-point semidefiniteness, unit-modulus roots, and optimal LP values are not used to establish the mathematical theorems. The analytic proof of attainment is Theorem I.52.2 together with Lemma I.53.1. The exact rational rounding checks verify the arithmetic inequalities in the recurrence conversion, not the external Weyl estimate. No Lean formalization is included or claimed.

V.138.2 Suggested integration boundary

A suitable new directory is

`Combinatorics/Ramsey/Research/GowersSzemerédi/sharp-missing-intervals/`.

The manuscript should initially remain a research source. Existing `Section05.lean` targets should not be silently redefined, and the status ledger should not be marked proved on the basis of this PDF. A separate collection of strengthened statements can preserve compatibility with the already formalized statements.

A practical formalization can be split into two tracks. The lower-bound track needs only the coefficient formula, positivity, a finite Fourier expansion, and summation against nonnegative weights. It does not need the moment-existence theorem, unit-circle root geometry, or measure separation. The equality track adds the supported moment criterion and the rank argument. This separation prevents the sharpness proof from becoming a dependency of the useful finite-cyclic inequality.

Proposed statement	Main mathematical dependencies
Positive gap certificate	Trigonometric recurrence; autocorrelation identity
Weighted linear-cutoff lemma	Certificate; finite Fourier sum; nonnegative weights
Exact first-window optimum	Arc positivity; Toeplitz/localizing matrices; rank
Noisy-gap uniqueness	Certificate slack; Vandermonde injectivity
Critical concentration	Sine polynomial; elementary sine lower bound
Recurrence conversion	Weighted lemma; Dirichlet; existing Weyl companion
Short-step interface	Recurrence conversion; integer square-root inequalities

V.138.3 A finite Fourier interface

The following is a mathematical specification, not Lean code:

For $N \geq 2$, $1 \leq M$, $2M \leq N$, and $w : \mathbb{Z}/N\mathbb{Z} \rightarrow \mathbb{R}_{\geq 0}$, if $\sum w > 0$ and w vanishes on $|x|_N < M$, there is an integer $1 \leq r < N$ with $r \leq N \operatorname{div} M - 1$ and

$$-\operatorname{Re} \hat{w}(r) \geq \frac{M}{N-M} \sum_x w(x).$$

Using an integer frequency and then casting it into $\mathbb{Z}/N\mathbb{Z}$ avoids ambiguity about centered representatives. A second lemma can convert to the existing `centeredAbs` interface. The empty-weight case should be handled separately rather than divide by zero.

For empirical measures, the implementation should use a finite sum over sample indices or a multiplicity function. Replacing it by the image set would change both the mass and the Fourier sum. The existing repository already contains the needed collision-safe pattern.

V.139 Further research problems

Source 37, delivered Section 11.

The following problems separate plausible extensions of the present calculation from much larger conjectural consequences. They are proposed research questions, not claims that the wider literature has no answers.

Question V.139.1 (Beyond the first window). *Determine $\inf_{\mu \in \mathcal{P}(K_a)} \max_{1 \leq r \leq R} |\hat{\mu}(r)|$ when $a > 2\pi/(R+1)$, including exact extremizers and transition points. Does a finite sequence of contact-pattern changes lead to an explicit piecewise algebraic description in $\cos a$ for each fixed R ?*

The prescribed equimoment matrix ceases to be positive once $\rho > 1/R$, so continuing the tangent formula is not justified. A successful extension must change either the active moment pattern or the primal support. The small-dimensional cases can be explored by the supplied LP code, but attainment and optimality still require analytic certificates.

Question V.139.2 (Exact finite-grid corrections). *For fixed N, M, R , solve the same extremal problem with support restricted to $\mathbb{Z}/N\mathbb{Z}$, and distinguish open from half-open exclusions. Quantify the excess over the continuous optimum in terms of the distances from the continuous contact nodes to the grid.*

The lower certificate survives sampling exactly, whereas the equality measure may not fit the grid. A second-order upper approximation is plausible when all relevant nodes can be moved within the allowed arc; endpoint constraints can produce a different first-order behavior. A complete result must account for both effects and positivity of the perturbed masses.

Question V.139.3 (Uniform stability constants). *Obtain useful bounds for κ and σ in Theorem I.56.3 as $R \rightarrow \infty$ with $(R+1)a$ in a compact subinterval of $(\pi, 2\pi)$. Can a suitable rescaling of angular distance give constants independent of R ?*

The critical estimate (I.56.1) suggests geometric control may be much better than a generic Vandermonde estimate. Weight recovery is likely the more delicate part. One should distinguish transport toward an unweighted contact set from transport to the uniquely weighted extremizer.

Question V.139.4 (Unequal Fourier budgets). *Given $\varepsilon_1, \dots, \varepsilon_R$, determine the least possible $\mu(I_a)$ subject to $|\hat{\mu}(r)| \leq \varepsilon_r$. Characterize the region of budget vectors for which the certificate of Section I.51 remains optimal.*

The available lower bound is $d - 2 \sum f_r \varepsilon_r$. Equality would require a measure on the same contact set with prescribed negative real moments. This turns local optimality into a finite feasibility problem; its positivity constraints may describe a computable cone or polytope in a suitable parameterization.

Question V.139.5 (Several forbidden arcs). *Find sharp Fourier detection results for a union of missing arcs, and compare their cutoff requirements with the total excluded length and the largest individual gap. Which arrangements of arcs force simultaneous large coefficients at several frequencies?*

The one-arc sign factor $\cos \theta - \cos a$ is no longer sufficient. Products of sign factors give certificates, but their Fourier coefficient signs and optimal degrees become nontrivial. This connects directly with multi-band moment representations rather than with a routine rescaling of the single-gap argument.

Question V.139.6 (Parameter optimization in the recurrence conversion). *Use the exact frontier $\rho_R(2\pi\delta)$, rather than the simplified $\delta/(1-\delta)$ bound, and optimize the cutoff, Dirichlet scale, and constant allocations simultaneously. Determine the best constant obtainable from the fixed Weyl input and this one-step architecture.*

The exponent $1/(D+1)$ in the present conversion reflects the powers in that architecture. Improving its numerical constant is a different task from proving that no more efficient recurrence argument exists.

Question V.139.7 (Propagating the short-step estimate). *Rework the degree induction for polynomial phase partitions using (I.57.16). What improved target-length exponent and threshold follow once every lower-coefficient term and every integer rounding loss is retained? Then propagate only the verified improvements through the later inverse and density-increment steps.*

This is the direct route from the present local work toward a stronger quantitative Szemerédi statement. A plausible local improvement is not itself the global theorem; a full parameter ledger is needed.

Question V.139.8 (Additional structure in polynomial empirical measures). *Can the moment problem be improved when the measure is known to be $T^{-1} \sum_{n \leq T} \delta_{an^k}$, rather than arbitrary? In particular, can its near equality with $\mu_{R,a}$ be excluded or exploited to force a better recurrence estimate?*

The extremizers here have $R+1$ specially placed atoms and negative equal moments. A polynomial orbit may approximate such a measure only for special arithmetic reasons. Turning that observation into a quantitative dichotomy would use the equality and stability theory rather than only the numerical coefficient bound.

Question V.139.9 (Certified symbolic extremizers and formalization). *For algebraic $\cos a$, produce exact algebraic descriptions of the contact nodes and positive weights, with certificates small enough for a proof assistant. Formalize first the finite lower-bound track and then the independent sharpness track described in Section V.138.*

The recurrence proof for w , the coefficient identity, and the rank-one matrix decomposition are natural certificate objects. Root-location verification need not be a separate numerical task because the moment-rank argument proves it structurally.

Question V.139.10 (Classical extremal interpretations and priority). *Identify the precise relationship of the formula (I.50.4) and the extremizing quadrature to classical trigonometric moment bounds, extremal positive polynomials, and finite power-sum inequalities. Which parts are known in equivalent language, and which refinements or stability statements are genuinely new?*

The general positivity machinery is already classical. Establishing the historical status of the closed formula and its consequences requires a more targeted literature comparison than recognizing that machinery. The proofs in this manuscript do not depend on a novelty claim.

V.140 Conclusion

Source 37, delivered Section 12.

The missing-interval step admits a sharp, exactly attainable Fourier formulation in the first spectral window. Its certificate has three useful features at once: the correct sign on the complementary arc, strictly positive nonconstant Fourier coefficients, and an explicit equality measure certified by a rank-one localizing matrix. This yields a robust mass tradeoff and rigidity, not just an isolated large-coefficient lemma.

For the selected Gowers argument, the immediate substitution is concrete: frequency at most $\lfloor N/M \rfloor - 1$ and coefficient at least $|A|M/(N - M)$, instead of the proof-supported quadratic cutoff and $|A|M/(4N)$. Optimizing the following rational approximation gives the explicit exponent b_k in (I.57.13), with its constant and short-step threshold displayed. The package is intended as a research source for careful integration and formalization. Global parameter propagation and historical priority remain separate tasks.

V.141 Source audit and reproducibility notes

Source 37, delivered Appendix C.

The following distinctions are material to integration.

Source or assertion	Treatment in this manuscript
Original Gowers Lemma 5.2	The proof's $1/4$ coefficient is the comparison baseline; the printed $1/2$ is not called false.
<code>Section05.lean</code>	Source for the corrected target, Fourier convention, and explicit thresholds.
<code>Proofs05MissingInterval.lean</code>	Existing proof-supported missing-interval development; not modified.
<code>Proofs05PolynomialPartition.lean</code>	Existing collision-safe multiplicity correction and integer-length issues are retained and credited.
<code>Proofs05Weyl.lean</code>	Its displayed <code>lemma_5_3_holds</code> supplies the cited explicit Weyl input; no build or re-verification is claimed here.
Supported moment criterion	Classical machinery, proved here for completeness and related to Yang–Xie.
Exact tangent formula and extremizers	Established by the proofs in this manuscript; historical novelty not asserted as settled.
Recurrence bounds	Consequences of a stated external Weyl estimate; no record comparison or global Szemerédi consequence asserted.
Computational checks	Finite symbolic and numerical regression tests, clearly separated from the proofs.

Build the document with `latexmk-pdfsharp_missing_intervals.tex`, or run `pdflatex` three times. The bibliography is embedded, so no external bibliographic database is needed. Run `pythonverify.py` after installing the packages in `requirements.txt`. This overwrites the local `verification_results.json` with a fresh report. Neither the build nor the verification script accesses the network or modifies a repository.

V.142 What the calculations verify

Source 38, delivered Section 8.

Delivered part heading of source 38: **Verification, integration, and further research**

The universal assertions in Parts I–III follow from the written proofs. The accompanying programs check finite identities, rational constants, symbolic series, and representative numerical instances independently. Their roles differ, and the output files identify which operations are exact and which use floating-point arithmetic.

V.142.1 Algebraic and finite checks

The exact symbolic calculation in `code/verify_quotient.py` reconstructs the fixed-phase norm solution and checks the coefficients through tenth order. Its separate finite-group calculation evaluates the full Fourier-square sum on $\mathbb{F}_5^2$, compares it with the centered polynomial, and checks the retained projection defect. It uses 160 reproducible random inputs, including functions close to a principal Fourier pair. Among these, 62 satisfy the local coercivity hypotheses with a numerically

resolvable positive norm loss. The largest centered polynomial discrepancy is less than $5 \cdot 10^{-16}$; the smallest observed improvement divided by $t^4 \Delta$ is greater than 3.18. These are consistency checks on examples, not estimates used in the dimension-uniform proof.

The separate program `code/verify_twelfth_order.py` checks the norm recursion, the Hessian, amplitude and phase displacements, and all terms in the optimized twelfth-order coefficient by exact symbolic arithmetic. It also verifies the signed-parameter symmetry. Its high-precision calculation solves the exact stationary equations at three positive parameter values. The mathematical compactness and implicit-function argument, rather than this numerical solver, establishes that the stationary branch gives every global maximum near zero.

For the broader odd-order expansion, the program `code/verify_odd_cube_expansion.py` directly enumerates physical cubes on $C_3, C_3^2, C_5, C_7, C_9, C_{15}$ and compares every coefficient with the Fourier formulas. Thus the test includes nonprime cyclic orders and order-three torsion. It is independent of the Fourier-square implementation in the characteristic-five test.

The standard-library program `code/verify_odd_rigidity.py` checks the exact three-point and five-point support multiplicities, then verifies the Weyl expansion and the finite envelope on 120 reproducible random vectors over $C_3, C_5, C_7, C_9, C_{15}, C_{25}$. The largest expansion residual is below $8 \cdot 10^{-16}$. Finally, `code/verify_six_ap_defect.py` performs the exact finite-field computations detailed in Part III. Those checks use integer arithmetic and exhaust all relevant field elements and target values for the two stated base fields.

V.142.2 High-precision illustrations

Figure V.142.1 displays two explicit families after subtracting the earlier terms in their asymptotic expansions. On the left, the five-point family uses the limiting phases and the corrected amplitude

$$b = (q_0/3)t^2 + (31q_0^2/27)t^4,$$

with the principal amplitude determined by the exact norm equation. This family attains the sharp tenth-order coefficient. Its phases have not been adjusted, so it does not attain the sharp twelfth-order coefficient. On the right is the exact three-point family in Part II. The variable ε there is its actual deficit $1 - \kappa$. The displayed residuals are

$$\begin{aligned} C_{10}(t) &= \frac{Q_{\text{fam}}(t) - 1 - 6\sqrt{2}t^4 - (2^{3/4}/3)t^6 - (20/9)t^8}{t^{10}}, \\ \mathcal{R}_{\text{odd}}(\varepsilon) &= \frac{t - \varepsilon/4 - 3\varepsilon^2/16}{\varepsilon^{5/2}}. \end{aligned}$$

The program `code/make_figures.py` writes the figure and its underlying CSV data. It uses 121 parameter values in each panel and records the precision and norm-equation residual. At the smallest plotted five-point parameter, the tenth-order residual differs from its limit by less than $4 \cdot 10^{-12}$. The high precision is useful because subtracting the preceding asymptotic terms at ordinary machine precision would conceal the coefficient being plotted.

V.143 Integration with the ProveIt development

Source 38, delivered Section 9.




Figure V.142.1: Residuals for two explicit families. The dashed lines are the proved limits $869/(216 \cdot 2^{3/4})$ and $\sqrt{2}/64$. The plotted values are computed at 120 decimal digits; they are illustrations, not numerical certificates of global optimality.

V.143.1 The exact connection to Gowers’ lemmas

The direct strengthening of Gowers’ Lemma 3.10 is Corollary III.85.5: for degree-two uniformity parameter α , it replaces the triangle-inequality envelope $(\delta + \alpha^{1/8})^8$ by the exact small-norm profile in the specified exponent-five setting. The fixed-odd-exponent theorem extends the structural reduction, even where the resulting cyclic profile has not yet been explicitly computed. These statements have narrower arithmetic hypotheses than the original all-modulus lemma and substantially more precise conclusions within that scope.

The proofs also strengthen how Lemmas 3.8 and 3.9 are used. Their usual upper estimates discard the location of Fourier mass. The retained energy inequality (III.84.1) and the linear seventh-degree estimate (III.84.4) turn a non-strict norm comparison into a strict improvement whenever a candidate has Fourier mass outside the selected cyclic subgroup. Part II addresses a different endpoint: it quantifies near equality in the critical U^3 – L^2 inequality, while Part III supplies an exact count for an extension of the layered norm method associated with the progression questions in Gowers’ Section 4.

V.143.2 Formal status and normalization

The inspected checkpoint `FORMALIZATION_STATUS.txt` reports 206 compiled modules, 1,878 audited public Gowers theorems, 100 exact companion proofs, and 20 open catalogue statements [168]. These are the repository’s reported results at the pinned commit. This project inspected the relevant source files; it did not reproduce that build. No new theorem in this article has been checked by the Lean kernel.

[write] **6 October 2026; sources 34–39.** *Checked at HEAD 89c1cbc38.* The counts above are those of the pin; at HEAD `FORMALIZATION_STATUS.txt` reports 410 modules and 2,840 audited public theorems, and the ledger 107 exact companions and 13 open statements (the source’s `SOURCE_LEDGER.json`, shipped as

`data/38-quotients-SOURCE_LEDGER.json`, keeps the pin’s 206, 1,878, 100 and 20; recorded, not edited). `lemma_3_10` and its companion are unchanged; the four Lean files named in the table below exist with the blobs the source records. No statement of this source is formalized.

The existing declaration `lemma_3_10` and its companion `lemma_3_10_holds` should retain their present scope. The new theorems are separate refinements with explicit realness, centering, odd-order or exponent assumptions. In particular, a proof on $\mathbb{F}_5^n$ cannot be substituted for a theorem on arbitrary $\mathbb{Z}/N\mathbb{Z}$ without changing its statement.

The repository’s Section 3 cube form uses unnormalized sums. For a group of order N , its three-cube norm is $N^{1/2}$ times the normalized U^3 norm used here, and its three-cube count is $N^4\mathcal{C}_3(W)$. A formal interface should prove these normalization equalities first. For general cube dimension d , the corresponding norm factor is $N^{(d+1)/2^d}$. Losing these factors would change both the hypotheses and the powers of the density in the result.

Mathematical component	Existing interface inspected	New proof work
Cube Cauchy–Schwarz and triangle inequality	<code>Proofs03GowersCS.lean</code> ; <code>Proofs03Minkowski.lean</code>	Transport to normalized finite-group averages and subgroup projections
Sharper cube upper bound	<code>Sections01_03.lean</code> ; <code>Proofs03CubeUpper.lean</code>	Exact Fourier coefficients, coercivity, quotient pullback, finite algebra
Odd-order critical stability	No matching theorem was identified in the inspected files	Weyl orthogonality, finite spectral projectors, quantitative local expansion
Six-term norm extensions	The research reports supply the preceding mathematical context	Norm fibers, trace quadratic forms, centered interpolation, exact counting

A practical order is to formalize the finite algebraic statements before the analytic extremal expansion. The Fourier-square identity, the rational cutoff, the six-point interpolation matrix, and the quadratic-form count have exact finite certificates. The analytic profile additionally needs compactness, an implicit-function theorem, and a justified Taylor remainder; symbolic output alone is not a formal proof of these steps. The programs are designed to expose the identities such a development would need, without presenting uncompiled Lean placeholders as checked results.

The archive supplies a complete TeX manuscript, its bibliography and figure, verification programs, and build instructions. The source ledger records the inspected repository checkpoint and preceding reports, preserving the distinction between earlier and new results.

V.144 Proposed further research questions and topics

Source 38, delivered Section 10.

The following questions are proposed continuations of the proved results. They are not claimed to constitute an exhaustive survey of open problems. The characteristic-five eighth-order question and the odd-order asymptotic coefficient question from the preceding reports should be replaced by the resolved statements above. The problems below concern the stronger conclusions that remain out of reach of the present arguments.

1. **Determine the sharp range of exact quotient reduction.** Define c_5^* as the supremum of c for which every exponent-five maximizer factors through a five-point quotient whenever $0 < u/\delta < c$. Is c_5^* finite, and what first destroys the quotient conclusion? The proof gives $c_5^* \geq 10^{-5}$. A useful next step is to compute the exact second variation against off-quotient perturbations along the analytic five-point branch, then distinguish a local instability from a crossing with a distant competing maximum.
2. **Decide whether the odd-exponent cutoff must depend on the exponent.** The sufficient constant $10^{-8}m^{-8}$ comes from summing bounds over up to m^2 Fourier-square slots. Can a global square-sum estimate remove that dependence, giving one positive cutoff for all odd exponents? If not, construct an explicit sequence of odd-order groups and genuinely noncyclic extremizers at parameters tending to zero. Either outcome would explain the precise scope of the projection principle.
3. **Classify the cyclic profiles and their arithmetic transitions.** Compute the first unresolved terms of $\mathcal{M}_{1,\mathbb{Z}/m\mathbb{Z}}(t)$ for small odd m , and determine which character order supports a maximum when m has several odd prime divisors. Order three already supplies a fifth-degree contribution, whereas the exponent-five profile has the secondary scale studied here. A systematic classification should record the frequency relations responsible for each power and decide when fractional powers occur.
4. **Find an exact algebraic description of the five-point branch.** The norm constraint and stationary equations give a finite algebraic system after replacing phase angles by their sines and cosines. Can it be eliminated to an explicit equation for the extremal value with a certified branch-selection interval? Determine the first nonzero term, if any, in $\alpha(t) - \pi/10$ after its third-order cancellation, and the symmetries behind subsequent phase coefficients. An answer would make the convergence radius and the higher coefficient recursion more concrete.
5. **Optimize the quotient-stability constants and powers.** Determine the best constants, and the sharp dependence on the extremal deficit, in (III.84.6)–(III.84.7). The factor 256 comes from a general convexity estimate and has not been optimized for this Fourier geometry. Constructing perturbations in one additional frequency coset is a controlled way to test which exponents and constants can be improved.
6. **Describe the constrained weighted profile beyond the interior regime.** When the maximizing five-point function first reaches 0 or 1, which coordinates saturate, and does exact quotient reduction continue? Determine the resulting piecewise extremal profile at a fixed density. Separately, quantify how fast indicator sets in growing product groups can approximate these weights in the uniformity norms relevant to the cube count. The variance in L^2 must be treated on its own terms.
7. **Determine the next sharp odd-order stability coefficient.** Find the coefficient of ε^3 in $\mathfrak{S}_{\text{odd}}(\varepsilon)$. Does the three-point family remain extremal at that order, or can perturbations outside its two tangent eigencharacters improve the value? The tangent-rigidity corollary narrows the search to perturbations around an order-three pair. Resolving the next variation would also indicate whether there is an exact finite-parameter odd-order modulus.
8. **Solve the coprime-to-six local extremal problem more precisely.** With order-three resonance excluded, determine the optimal cubic term in the deficit and whether a fixed cyclic group realizes it. The two-state five-point family proves the displayed second-order sharpness, but it need not settle the subsequent coefficient. This is a separate problem from optimizing the coarse constant 104 in the finite envelope.
9. **Make local rigidity constructive from sampled derivative data.** The proof identifies a large Weyl subgroup and then a common eigenvector. Can this argument be converted into an efficient reconstruction procedure with error constants approaching the sharp modulus? The relevant input may be an approximate ambiguity table or samples of it, rather than the complete function. A useful formulation should distinguish arithmetic cost, sampling cost, and the dependence on the exponent.

10. **Develop sharp critical stability for higher Gowers norms.** For U^d , the natural critical exponent is $p_d = 2^d/(d+1)$ [7]. Determine the leading stability coefficient near a polynomial phase supported on a coset, and identify any torsion-dependent fractional corrections. The exact L^2 Weyl identity used here is specific to $d = 3$; a higher-order argument must replace that Hilbert-space mechanism.
11. **Calculate the eight-term and higher norm-extension defects.** The next complete core contains norms of degrees 1, 3, 5, 7. Determine the joint distributions of the polarized coefficients that remain after the top direction is fixed. In the six-term case only the quadratic statistic e_2 was needed; higher layers bring several interacting coefficients. An exact formula, or a uniform error estimate around its natural mean, would extend Part III in a substantive way.
12. **Obtain a stability theorem for small norm-extension defect.** If $A_6^{\text{sym}}(P)$ is close to its baseline, how close must R be to an additive quadratic? Formula (III.88.1) first controls the fraction of nonzero third and fifth differences. The remaining task is a quantitative polynomiality theorem with constants uniform in the number of added coordinates. Determine the correct dependence on the characteristic and on the extension degree of the base field.
13. **Resolve the global quintic four-term avoidance dimension.** The preceding degree-stratification report leaves the maximum dimension for a degree-five coloring without a nontrivial symmetric four-term progression between four and eight [173]. The independent-extension obstruction does not resolve this gap. Can a genuinely coupled quintic construction exceed the known four-dimensional construction, or does a stronger global argument force a smaller upper threshold than the pointwise degree count?
14. **Certify the new exact identities in the formal development.** Begin with the centered cube coefficients and the energy coercivity, then the finite-field fiber calculation and the Weyl-coset expansion. These pieces can be reviewed separately and are useful without the analytic optimization. A later proof of exact quotient reduction should preserve the strict deficit and the explicit cutoff, rather than formalizing only a weakened asymptotic bound.
15. **Locate a genuine bridge to global progression bounds.** Identify a density-increment or transference argument in which one of these sharper local statements survives all subsequent losses. Such an application must explicitly reconcile the bounded-exponent setting with the ambient group used to model long integer progressions. Write out the resulting recurrence and solve it before claiming an improvement of the global bound. This problem is logically additional to the local extremal theorems proved in this article.

V.145 A precise interface with the inspected Gowers development

Source 39, delivered Section 12.

V.145.1 Source controls and the already existing improvements

In this section q denotes the number of vertical affine classes and b the complexity parameter of each horizontal section. We deliberately do not use s for both a class size and a section complexity.

Let $c(u, \gamma, k)$ and $Q(u, \gamma, k)$ denote the source's positive scale and graph controls, represented by `multipleC` and `multipleQ`. For a horizontal section of complexity b , the relevant one-slice form is

$$a(\tau) = c(\tau/b, \gamma, k)^b, \quad L(\tau) = Q(\tau/b, \gamma, k)^b. \quad (\text{V.145.1})$$

These symbols retain their source meaning; the present article does not redefine their dimension-dependent exponents.

The inspected `Proofs16RecoveredSliceCover.lean` already gives a common refinement with

$$P \leq Q(\epsilon/(rb), \gamma, k)^{rb}, \quad (\text{V.145.2})$$

$$\text{base width} \geq M^{c(\epsilon/(rb), \gamma, k)^{rb}}, \quad (\text{V.145.3})$$

and a recovered candidate index set of cardinality

$$rP + r^2P^2. \quad (\text{V.145.4})$$

Its geometric successor imposes a short-parent width condition and performs synchronization without a further candidate blow-up [177, 179]. In particular, the linear dependence on r in the refinement exponent is *already present* in this snapshot; it is not a new conclusion being attributed to this article.

V.145.2 Replacing the sample while keeping the existing interface

Proposition V.145.1 (Sampled-or-anchored output). *Suppose ϕ on a partial subset D of $H \times J$ is covered, on every vertical fibre, by at most q affine functions over a field. For every $1 \leq r \leq m$, there is an injectively enumerated sample of r columns and sets $K \subseteq F \subseteq H \times J$ such that:*

- (i) *the complement of F has ambient mass at most $2q(m-r)/(m(r+1))$;*
- (ii) *every point of $D \cap F$ is either in K or on a sampled column;*
- (iii) *every point of $D \cap K$ has two distinct sampled anchors in the same row, both in $D \cap K$, from whose values it is recovered by (II.159.4).*

These are the mathematical contents needed by the repository's `Section16SampledOrAnchoredOn` interface.

Proof. In each row assign each point of D to one of the affine graphs that covers it, for example the first one in a fixed ordering of the q graphs. This makes a disjoint class partition. Apply Theorem II.153.2 with $t = 2$. Let $F = E_R$ be its ambient good set and let K be the union, over all rows, of the full classes containing at least two sampled columns.

Every such full class is contained in F , so $K \subseteq F$. A domain point of F not in K survives precisely because its column is sampled. For a point of K , choose two sampled members of its own class. They too belong to that full class and hence to $D \cap K$. They have distinct column values, and the affine interpolation identity follows from the common affine class function. Enumerating the set R produces an injective sample. $\square$

[write] 6 October 2026; sources 34–39. Checked at HEAD 89c1cbc38.

Section16SampledOrAnchoredOn is at Proofs16ShortColumns.lean, line 27; the seven Proofs16 files the source cites have the blobs it records. The proposition is a formalization candidate (an exact double count over $(r+1)$ -subsets), claimed by nobody; the same $O(q/\sigma)$ sampler is source 19's (note in Section F.44), which should be credited first. lemma_16_10 (Section16.lean, line 851) is open; nothing here changes that, as the source says.

The point that the anchors belong to K , not just to the larger set F , is useful: it matches the nested anchored-subset predicate in the repository. This is why we take K to contain whole sufficiently sampled classes, rather than just the nonsampled targets.

In the finite-field implementation $J \subseteq \mathbb{Z}/N\mathbb{Z}$ with N prime, distinct column values have invertible difference. The set-theoretic sample theorem itself does not need primality. Over a general ring the same algebraic interface would need unit differences for the chosen anchors; nonzero difference alone would not suffice.

V.145.3 A parameter-explicit local replacement

Theorem V.145.2 (Local affine lift with the existing common-cover bounds). *Assume the vertical affine cover of Proposition V.145.1, and assume the horizontal sections satisfy the source's common-cover hypothesis with parameters (γ, b, k) . For $0 < \eta, \epsilon < 1$, choose*

$$r = \min\{m, \lceil 2q/\eta \rceil\}. \quad (\text{V.145.5})$$

At the level of base cells times the final-coordinate set, the existing common-cover and interpolation conclusions apply with exceptional fraction at most $\eta + \epsilon$, width bound (V.145.3), and at most $rP + r^2P^2$ candidates, with P as in (V.145.2).

If, in addition, the labelled hereditary hypothesis (V.145.1) is used while preserving the required admissible base cells, the same width bound is obtained with candidate bound

$$rL + \binom{r}{2}L^2, \quad L = Q(\epsilon/(rb), \gamma, k)^b. \quad (\text{V.145.6})$$

Proof. Proposition V.145.1 and Corollary II.153.3 provide the sampled-or-anchored good set with error at most η and an r -term sample. The already established common-cover and recovered-slice conclusions then give the first assertion, losing another ϵ in the base. Equivalently, their purely algebraic covering step is the $t = 2$ case of Theorem II.160.1, with the common list bound P used for every anchor and an ordered-pair indexing convention.

For the labelled version, apply Theorem II.161.2 with $\tau_i = \epsilon/r$ and the one-slice controls (V.145.1). The product of the scale powers is

$$\prod_{i=1}^r c(\epsilon/(rb), \gamma, k)^b = c(\epsilon/(rb), \gamma, k)^{rb},$$

whereas each individual list has size at most $Q(\epsilon/(rb), \gamma, k)^b = L$. Corollary II.160.2 supplies (V.145.6). The same error addition applies. $\square$

The first part is a direct mathematical replacement at an existing interface; it requires no redesign of geometric synchronization. The second part is a proved refinement theorem with explicit hypotheses and a proposed way to retain more information in the implementation. It is not claimed that the new labelled version has already been formalized in the repository.

If the maximum permitted common list size is written as $P = L^r$, the old upper-budget expression is $rL^r + r^2L^{2r}$, while (V.145.6) is $rL + \binom{r}{2}L^2$. This comparison concerns *guaranteed list budgets*, not necessarily the actual number of distinct graphs produced by either method. Both arguments may have extensive redundancies in a particular instance.

V.145.4 Geometry must remain a separate hypothesis

A product of a refined base box and an arbitrary final-coordinate progression is not automatically a proper common-step box. The repository explicitly separates these issues. Its synchronization statement [178] uses a short parent progression for one axis of each refined base cell, parallel to the original final progression I , and with unit step. The parent is short relative to the ambient modulus, has length at most $|I|$, and contains the chosen refined axis. A further inequality of the form

$$v^2 \leq \text{refined base width} - 1$$

allows compatible final-axis tiling. Under these hypotheses it produces proper common-step boxes of width at least $v - 1$, without increasing the list of candidates or the exceptional set.

In the source’s assembled short-parent version, a sufficient real-valued budget is

$$v^2 + 1 \leq (m_0/8)^{c(\epsilon/(rb), \gamma, k)^{rb}}, \quad (\text{V.145.7})$$

where $m_0 \geq 4$ is an input common lower width, not the column-set cardinality m used in our sampling lemmas [179]. Appropriate properness, unit-step, and containment assumptions are also required.

Corollary V.145.3 (Conditional synchronized lift). *If the base refinement supplied in Theorem V.145.2 has the short-parent containment and width hypotheses of the source synchronization lemma, its rectangular cover can be replaced by the corresponding proper common-step box cover with the same exceptional fraction and the same candidate bound. This applies to the labelled bound (V.145.6) when the labelled refinement preserves those hypotheses.*

Proof. Synchronization partitions each rectangular product into compatible boxes and restricts the candidates from its base cell to the resulting boxes. Restriction preserves the candidate identities and multilinearity. The source tiling covers the whole rectangle, so there is no additional discarded mass. Its width conclusion is exactly the assumed synchronization conclusion. $\square$

This corollary is conditional on stated geometric input. It must not be read as proving that any abstract common refinement automatically satisfies that input, or as replacing an unproved scale comparison by the word “synchronization.”

V.145.5 Equal-error budget comparison

For a final sampling error allowance η , the inspected corrected sampler can use $\tau = \eta/2$ and the sufficient budget

$$r_{\text{corrected}} = \left\lceil \frac{24q}{\eta^2} \right\rceil, \quad (\text{V.145.8})$$

with the short-column alternative handled as in the source [175, 176]. Our sufficient distinct-anchor budget is (V.145.5). The comparison is at the same error allowance; comparing $6q/\tau^2$ with $2q/\eta$ without noting $\eta = 2\tau$ would conceal a change of convention.

As a numerical illustration, take $q = 100$, $\eta = 0.01$, and a column set large enough that neither of these budgets needs truncation. Formula (V.145.8) gives 24,000,000, while (V.145.5) gives 20,000. This is a comparison of sufficient local budgets, not an empirical running-time measurement. The asymptotic minimax budget in a sufficiently sparse-sampling limit is about 8,399.62, but that number is *not* a certified finite replacement for 20,000 at an unspecified value of m .

V.145.6 The closing comparison is not proved here

The inspected file `Proofs16ClosingComparison.lean` proves that the same-dimensional numerical comparison in the edited closing argument has the wrong direction [180]. In its parameter range,

$$Q(\rho, \gamma, k) < Q\left(\frac{\rho}{4p}, \gamma, k\right)^p \quad (0 < \gamma, \rho \leq 1, p \geq 1). \quad (\text{V.145.9})$$

It also records an endpoint obstruction with possibly different source and target dimensions: $Q(1, 1, d) = 1$, while shrinking the positive error argument below 1 and raising the corresponding graph budget to a positive power produces a number greater than 1.

The new sampler changes r , and the labelled refinement changes the local candidate count. Neither fact reverses (V.145.9) or proves every prescribed dimension-dependent control. Theorem V.145.2 therefore states the output bounds actually obtained. We do not conclude a new global bound for Szemerédi’s theorem, do not declare Lemma 16.10 fully repaired, and do not treat failure of a displayed proof step as a counterexample to the full mathematical lemma.

V.146 Exact checks, examples, and reproducibility

Source 39, delivered Section 13.

V.146.1 What was actually tested

The package contains a dependency-free Python implementation of the finite formulas, the dynamic program, the closed balanced-profile formula, the affine quadratic cap, and the deterministic sampling algorithm. The test programs use Python integers, integer square root, and `fractions.Fraction`; they do not use floating-point arithmetic for any assertion of equality or inequality.

Two certificate files record a total of 11,212 checked cases, all passing in the delivered run. Their categories are listed in Table V.146.1. A “case” means one loop instance in the named category; it is not a separately proved theorem or an independent review. Some cases contain several assertions. In particular, the second-difference cases compare the direct difference, the hypergeometric formula, and, when applicable, the factored sign formula.

Table V.146.1: Executed exact-arithmetic verification categories.

Category	Cases
Direct one-class averages and the $(r + 1)$ -sample identity	2,020
Dynamic programming versus exhaustive size profiles	315
Anchor sets in exact permutation-orbit extremizers	105
Weighted deterministic selection instances	80
Conditional-expectation identities	226
Nonsampled finite-field polynomial recoveries	257
Candidate index-set cardinalities	180
Closed balanced formula versus dynamic programming	2,496
One-class unimodality and initial discrete concavity	624
Exact second-difference identities	2,765
Affine quadratic cap versus exact marginal maximization	2,144
Total	11,212

The first certificate suite includes finite-field reconstruction over $\mathbb{F}_5$, $\mathbb{F}_7$, and $\mathbb{F}_{11}$, with thresholds 1 through 4. It checks Lagrange reconstruction at nonsampled targets having enough valid anchors; the already sampled targets require no interpolation. The second suite checks the closed finite formula against the independent dynamic program for $m \leq 13$, $q \leq 4$, thresholds through 6, and all sample sizes. The affine cap is compared against an exact scan through all class sizes for every $m \leq 64$ and $0 \leq r \leq m$.

These tests are designed to expose normalization errors, endpoint mistakes, off-by-one errors in distinct-hit thresholds, incorrect ceilings, and false maximizer claims. They supplement the proofs; none establishes a universal statement by finite exhaustion.

V.146.2 Representative exact finite values

Table V.146.2 shows values from `data/exact_profiles.csv`. The profiles listed there were originally found by the independent dynamic program. The closed balanced formula gives the same values, though in linear or plateau cases it may choose a different maximizing profile. This is expected: Theorem II.155.3 proves the existence of a balanced extremizer, not uniqueness of all extremizers.

Table V.146.2: Exact minimax values and simple universal bounds. All losses are fractions of ambient mass.

m	q	r	t	A maximizing profile	$\mathfrak{M}_{m,q,r,t}$	Bound (II.153.3)
20	2	2	2	(10, 10)	27/38	1
20	2	4	2	(7, 7)	1001/2850	16/25
20	2	8	2	(3, 3)	143/950	4/15
20	2	12	2	(2, 2)	2/25	8/65
20	2	16	2	(2, 2)	1/25	4/85
40	3	2	2	(13, 13, 14)	43/50	1
40	3	4	2	(13, 13, 13)	3159/5624	1
40	3	8	2	(7, 7, 7)	56637/209950	8/15
40	3	8	3	(10, 10, 10)	233595/525844	4/5
40	3	12	3	(7, 7, 7)	6867/25160	63/130

The table makes visible the distinction between a useful easy bound and the actual optimum. It also shows why equal-size classes with unused columns cannot be ruled out merely by assuming that a difficult example should use the whole domain.

V.146.3 Reproducing the results

From the package directory, run

```
python3 code/verify.py
python3 code/verify_balanced.py
python3 code/poisson_constants.py
make pdf
```

The first two commands regenerate the exact JSON certificates. The constants script uses high-precision decimal arithmetic only for the illustrative Poisson constants, clearly separated from the exact combinatorial checks. The PDF can alternatively be built by running `pdflatex` twice on `article.tex`. The manuscript uses no external figure files and no bibliography processor.

The code deliberately keeps both the dynamic program and the closed formula. Agreement between implementations using different mathematics is a useful regression check, although it is not a substitute for the proofs. The repository interface is documented in `FORMALIZATION_PLAN.md`; that file is a plan, not an executable Lean proof.

V.147 Further research questions and proposed next steps

Source 39, delivered Section 14.

The finite minimax value and the existence of balanced-and-capped extremizers are settled by the theorems above. The following are distinct follow-on questions; none is used as a hidden premise of the proved statements.

Research question V.147.1 (Classification of all finite extremizers). Characterize every maximizing profile, not just the balanced one, and every row family attaining the minimax value. The proof identifies equalization steps that preserve or improve the objective. Equality should be controlled by zero discrete curvature, flat first differences, and the distribution of rowwise losses across samples. A useful result would give a complete criterion for equality in Theorems II.154.2 and II.155.3, including the many nonunique cases with large sampling fraction.

Research question V.147.2 (Finite corrections to the golden-ratio law). Obtain explicit uniform error bounds for (II.157.7), separating finite-population corrections, integer rounding, and the approach of q/r to the transition. The affine radical formula makes the optimizing class size explicit, so an expansion of the hypergeometric marginal at that size could yield certified second-order sampling budgets rather than only an $o(1)$ term. Near the phase transition, rounding and whether the column budget is active should be treated together.

Research question V.147.3 (General-threshold cap computation). For fixed $t > 2$, derive an equally transparent description of the first maximizing size $a_{m,r,t}$. The first-difference identity is

$$\Delta \ell_{m,r,t}(s) = \frac{m-r}{m^2} [\mathbb{P}\{Z < t\} - t\mathbb{P}\{Z = t\}], \quad Z \sim \text{Hyp}(m-1, s, r).$$

After factoring binomial terms this gives a bounded-degree algebraic expression when t is fixed. The problem is to isolate its relevant root with exact boundary handling and a complexity bound independent of a scan through all m sizes.

Research question V.147.4 (Thresholds growing with sample size). Determine the asymptotics of λ_t and κ_t as $t \rightarrow \infty$, and prove a uniform version of Theorem II.157.2 when t also grows with r and m . The fixed- t proof uses finitely many low-count probabilities and does not establish such a regime. This extension would quantify the cost of interpolating increasing polynomial degree rather than only a fixed degree.

Research question V.147.5 (Algebra-aware improvements beyond certification). Find hypotheses under which the permutation-orbit lower bound ceases to apply because additional algebraic information is available. Examples include classes known to share a polynomial, restrictions on how class labels vary with the base point, or compatibility among overlapping fibre interpolants. The objective should be a precisely defined stronger recovery rule with its own upper and lower bounds, not an unsupported claim that arbitrary polynomial reconstruction requires $\kappa_t q / \sigma$ anchors.

Research question V.147.6 (Compressing labelled candidate lists). When can the upper count $\sum_a L_a + \sum_{|S|=t} \prod_{a \in S} L_a$ be reduced by algebraic dependence among lists? Is there a useful lower-bound model in which its leading order is necessary? Even the affine case may benefit from canonicalizing pairs of slice candidates that give the same lifted graph. Any proposed compression must retain the pointwise existence guarantee without assuming a globally consistent choice of branches.

Research question V.147.7 (Formal label preservation and geometry). Formalize the new sampler into the existing `Section16SampledOrAnchoredOn` interface, then formalize labelled hereditary refinement with a dependent list type indexed by the sampled column. Prove that the short-parent cover construction preserves the needed containments for the labelled lists. This is a concrete integration task: its goal is the exact bounds of Theorem V.145.2 and Corollary V.145.3, not a silent change to the source definition of multiply-linear control.

Research question V.147.8 (Nonuniform column costs and sampling designs). Extend the finite game to weighted columns or nonuniform sampling costs. Uniform row weights cause no difficulty here, and arbitrary row weights are already covered; nonuniform *column* weights are different. The $(r+1)$ -subset identity and permutation-orbit minimax argument depend on the uniform column law. A replacement should specify the admissible sampling design, the target mass measure, and the correct notion of a profile.

Research question V.147.9 (Robust recovery when good classes are unknown). Develop an algorithm that exploits Theorem II.159.5 without being given the good subset or its class partition. The additive error bound shows that correct interpolants exist in the generated list; it does not identify them. A useful result would combine shared sampling with an explicit consistency test or a list-decoding guarantee, with running time stated in terms of the actual representation of the fibres.

Research question V.147.10 (Repairing the global control recurrence). Replace the invalid closing comparison by a valid dimension-dependent recurrence for scale, graph count, and exceptional mass. The new local budgets should be substituted before attempting domination by a closed formula. Endpoint-safe controls may need a multiplicative slack factor strictly larger than one, since a graph budget equal to one at $(\rho, \gamma) = (1, 1)$ cannot dominate every nontrivial construction cost. The task is then to propagate the repaired recurrence through the subsequent sections and explicitly compute what global quantitative statement follows. This remains separate from the finite sampling problem solved here.

V.148 Conclusion

Source 39, delivered Section 15.

The shared interpolation step admits a direct mass-weighted treatment. The finite identity (II.153.1) eliminates the need to discard small classes and then impose a separate failure probability on every surviving class. It yields a uniform $O(q/\sigma)$ sample budget, while exact minimax symmetrization proves that the associated certification problem has no hidden dependence on the number of rows.

The finite problem is solved more explicitly than an averaging bound alone would suggest. Hypergeometric finite differences show that a balanced profile, capped at the first one-class maximum, is optimal. For affine interpolation the cap is the ceiling of an explicit quadratic root. In a sparse-sampling limit this finite structure becomes the golden-ratio coefficient $\varphi^3 e^{-\varphi}$, and the full transition between partial and saturated column profiles is given by (II.157.8).

On the algebraic side, interpolation is naturally performed from labelled slice lists. Retaining those labels gives the candidate count $rL + \binom{r}{2}L^2$ rather than using a common union list at every anchor. The theorem is compatible with the repository's existing sampled-or-anchored and synchronization interfaces; subject to the explicitly retained geometric hypotheses. It improves local information and local parameters; it does not supply the missing global closing comparison.

V.149 Proof dependencies and formalization targets

Source 39, delivered Appendix A.

V.149.1 A dependency map

The finite sampling and minimax results are independent of fields, polynomials, Gowers norms, and the source's control functions. Their logical chain is

$$\text{finite double count} \implies \text{shared-sample bound} \implies \text{minimax upper bound},$$

with permutation symmetrization supplying the matching lower bound. The finite-difference identities then give one-class shape, followed by balancing and the closed finite formula. The affine cap uses only a quadratic sign computation after the first-difference identity.

The Poisson results use these finite formulas plus an elementary product inequality, elementary limits, and single-variable calculus. They are not needed to obtain the nonasymptotic $O(q/\sigma)$ improvement or to compute the finite optimum.

The algebraic lift uses interpolation and the already proved sample certificate. The labelled refinement theorem uses restriction closure, finite partitions, and measure addition. Only the final source-interface section uses Gowers's particular single-slice and geometric hypotheses. Separating these layers should make a formal development substantially easier to audit than a single theorem combining all of them.

V.149.2 Suggested declaration boundaries

A proof-assistant implementation can first state everything in finite types and unnormalized counts. The following names are proposed; they do not refer to declarations already present in this delivery:

sharedAnchorDiscardCount

The sum of $(|C| - |C \cap R|) \mathbf{1}_{\{|C \cap R| < t\}}$ over disjoint row classes.

sharedAnchor_doubleCount

The bijection between (R, x) pairs and (T, x) pairs with $|T| = r + 1$. Proving the unnormalized integer identity first avoids probability-library and real-division overhead.

exists_sharedAnchorSet_mass

The finite average argument and the bound from $Z \mathbf{1}_{1 \leq Z \leq t} \leq t$. This is the shortest route to the local $O(q/\sigma)$ improvement.

sharedAnchor_minimax_profile

The symmetric-group action and its equal-size fibres, followed by the exact profile maximum. This is useful for optimality but is not a dependency of the interpolation application.

sharedAnchor_loss_secondDifference

The two distinguished-column coupling identities and the sign factorization in (II.155.4).

sharedAnchor_balanced_profile

Clipping at the first mode, filling the budget, and discrete balancing. The finite affine root may be handled by squared inequalities and an integer ceiling characterization rather than numerical square roots.

section16_subset_recovered_good_set

An application of the new sample bound that produces exactly the existing sampled-or-anchored interface. The whole-class set K in Proposition V.145.1 ensures that the anchors also belong to K .

labelled_common_slice_cover

Inductive refinement carrying a separate candidate family for each slice. A dependent index type avoids padding every slice list to the union size.

A formalization should establish an actual build at a pinned Lean/mathlib version before changing the status of any source lemma. This package contains no `sorry`-based Lean skeleton, no declaration of the new results as axioms, and no claim that a proposed theorem name has been compiled.

V.150 Notation and endpoint checklist

Source 39, delivered Appendix B.

Symbol	Meaning and restrictions
m	Number of final-coordinate columns; always a positive integer.
r	Number of distinct sampled columns; $0 \leq r \leq m$.
t	Number of distinct anchors sufficient for a class; degree bound is $t - 1$.
$q, q_h, \bar{q}$	Uniform, rowwise, and weighted-mean class counts. The finite minimax formula fixes an integer $q \geq 1$.
B_h	Partial domain in row h ; it need not equal J .
$K_h(R)$	All sampled domain points together with every class having at least t sampled members.
E_R	Ambient good set obtained by deleting uncertified domain points only.
$\ell_{m,r,t}(s)$	Expected discarded mass of one size- s class, normalized by m .
$a_{m,r,t}$	Least size attaining the maximum of ℓ , not its largest maximizing size.
$\mathfrak{M}_{m,q,r,t}$	Supremum over row families of the best shared r -sample loss for the stated rule.

Symbol	Meaning and restrictions
λ_t, κ_t	Unique Poisson-profile maximizer and its value. These are not finite-population constants.
b	Horizontal-section complexity in the source comparison; distinct from class size s .
M, m_0	Initial base scale and source short-parent width input; neither denotes $ J $ unless specified.
L, P	Individual slice-list and common union-list bounds. Keeping them distinct is the point of labelled refinement.

For $r = m$, every sampling loss is zero and the first mode is zero. For $r < t$, only sampled targets are certified and the minimax loss is $1 - r/m$. For $t = 1$, the direct lifts duplicate some one-point interpolation expressions, which is harmless for an upper bound. The finite formula uses a first mode so that plateau cases have an unambiguous cap. Sampled columns are distinct by construction. Field interpolation never divides by a zero difference. Exceptional mass is always ambient mass, and the corruption and base-refinement losses are added explicitly rather than hidden in a changed normalization.

V.151 Reproducibility, implementation, and formalization

Source 40, delivered Section 10.

V.151.1 What the finite checks establish

The accompanying standard-library Python program `verify_fejer_selection.py` uses exact integer and rational arithmetic for its mathematical checks. The JSON report records the chosen points and exact survival masses for every displayed orthogonality example. Its checks are as follows.

Check	Coverage and interpretation
Top-monomial rigidity	All 6 651 coefficient vectors over $\mathbb{F}_3$ for $k = 1, 2, 3$, with sidelengths one.
Total arrangement count	All 512 subsets of $\mathbb{F}_3^2$ for $k = d = 1$, including empty and full sets.
Small-field exceptional logic	All 625 ambient arrangements over $\mathbb{F}_5$ for $k = d = 1$, $H = 1$; all are exceptional, consistent with the bound, which is vacuous at that field size.
Kernel inequalities	156 exact (s, L) cases, for $s \in \{4, 8, 16, 32\}$ and $2 \leq L \leq 40$.
Selection constants	27 exact choices of s, α, η , checking the gain, ceiling estimate, score margin, and explicit retained mass.
Exact orthogonality	Four regular examples over $\mathbb{F}_{1000003}$: $(k, d, L) = (1, 1, 2), (1, 1, 3), (2, 1, 2), (1, 2, 2)$; full enumeration of bounded coefficient vectors, and exact good/bad Fourier sums.

The regular examples use $\phi = F$ for the good case. To form a bad case, the value of ϕ at one distinct vertex is increased by one. This changes the defect by a nonzero signed unit. The program checks that the good moment is exactly $a_{s,L}$ and the bad moment exactly $b_{s,L}$, not merely close in floating-point arithmetic.

The tests do not enumerate the enormous ambient fields in the main theorem, do not prove the assertions for all parameters, and do not verify a Lean development. The proofs in the article, not the sample computations, are the basis for the general theorem.

V.151.2 An explicit finite search procedure

The existence argument admits a finite, though generally impractical, construction. Enumerate the q^2 seed pairs (t, u) and compute the conditional expectation of the score $X - Y/\eta$ by enumerating labeled arrangements. Some seed pair has conditional expected score at least the unconditional lower bound. For this pair, expose the Bernoulli membership decisions one physical point at a time. At each decision, one of the two conditional expectations is at least their weighted average, so choose that outcome. At the end the realized score is at least the original guaranteed score.

This is a conditional-expectation proof of finite constructibility, not a near-linear-time algorithm. Direct evaluation can involve the q^M arrangement parameter space and $|B|$ decisions. If exact comparison of the trigonometric weights is desired, they lie in a cyclotomic number field; a certified algebraic-real representation can be used. For arbitrary real input η , exact effective comparisons also require a representation of that input. Taking rational accuracy parameters avoids this additional computability issue. No fast implementation of this construction is claimed or supplied.

V.151.3 Suggested Lean dependency structure

The existing mathematical project supplies a natural home for a formalization, but this package does not change its Lean files. The following proposed components isolate the substantial obligations.

Proposed component	Mathematical content
<code>CubeTopMoment</code>	Subset-sum inversion and the parity criterion of Lemma II.166.1.
<code>ArrangementFeature</code>	Parameterized arrangements, elimination of the last cross-section, and Theorem II.166.2.
<code>ShortRelationCount</code>	Finite coefficient boxes, sign pairing, polynomial zero-count, and repeated-vertex exclusion.
<code>PositiveKernelMoment</code>	Finite character orthogonality, nonnegative Fourier sums, and actual-versus-labeled survival.
<code>FejerSelection</code>	The score lemma, the exact theorem, ceiling arithmetic, and the density corollary.
<code>KernelGainBarrier</code>	The nonnegative-Fourier gain-mass inequalities, independent of arrangement geometry.

These are suggested names, not assertions that such modules already exist. A formalization should first fix the intended arrangement counting convention, particularly repeated vertices, before adapting existing interfaces. No uncompiled proof sketches with placeholder axioms are included as if they were verification.

V.152 Further research questions and integration priorities

Source 40, delivered Section 11.

The questions below distinguish a direct follow-up computation from a structural conjecture. None is asserted solved by the preceding proofs.

Research question V.152.1 (Sharp finite-bandwidth kernel optimization). For fixed s, H, Q , maximize $\sum_{a=-H}^H c_a^s$ subject to $c_{-a} = c_a \geq 0$, $0 \leq \sum c_a e^{ia\theta} \leq 1$, and $\sum c_a^s \geq Qc_0^s$. How close is the Fejér kernel to the optimum, including its dependence on s ?

The exponent is determined by Section II.172, but the constant and the necessary bandwidth are not. A useful first target is a rigorous solution for small H , with certified positivity rather than a sampled grid of angles. This would clarify whether better kernels can materially reduce the field threshold even when they cannot improve the powers of α and η .

Research question V.152.2 (Counting distinct short-relation hypersurfaces). How much can one improve $C_{k,s,H}$ by identifying scalar-equivalent or coincident zero sets among the coefficient vectors, and by using the geometry of their intersections?

Sign pairing is only the first such identification. Over a prime field, a projective count sensitive to the actual integer coefficient box may save more than a constant. Any refinement must still include zero sidelengths and coincident vertices. A credible target is an explicit bound of the form q^{M-1} times a substantially smaller coefficient complexity, rather than an unproved independence heuristic.

Research question V.152.3 (Threshold versus intrinsic complexity). Is a field-size requirement of roughly $\beta^{-(2d-1)}(s/(\alpha\eta))^{2s}$ necessary for an unconditional one-stage theorem of this kind, or can it be reduced while retaining the same arrangement mass?

Our requirement comes from making all exceptional bad arrangements cost at most the ordinary bad first moment. It may be wasteful to bound each exceptional arrangement's survival by one. Stratifying extra-relation ranks and calculating their survival could improve the threshold without changing the kernel or its exponent barrier.

Research question V.152.4 (Small-characteristic selection). Can a different family of probability kernels recover comparably strong selection in characteristic $p \leq 2(L-1)$, with a quantitatively controlled dependence on p ?

Two issues must be separated: modular collisions in the feature coefficient box and low-order target defects. Proposition II.171.2 computes the latter exactly. A successful approach cannot simply increase the extension degree of the field and ignore its fixed characteristic. Vector-valued frequencies, structured zero sets, or several genuinely different tests are plausible directions, but are not proofs.

Research question V.152.5 (Multiple universal relations). For a parameterized configuration whose polynomial feature relations form a lattice or vector space of rank greater than one, can the single-stage method be developed with constants controlled by the geometry of that relation space?

The exact Fourier sum is still available, but the good contribution becomes a weighted sum over a higher-rank coefficient section. Different nonzero defects can annihilate different subspaces, so the one-dimensional ratio calculation must be replaced rather than merely renamed. This is a natural route toward broader multilinear and hypergraph configurations.

Research question V.152.6 (Propagation through the remainder of Gowers' proof). What explicit bounds result from substituting Corollary II.170.1 into the subsequent strong-partial-multihomomorphism and multilinear-structure steps, while tracking every density, accuracy, and largeness condition?

The immediate task is a parameter ledger for each later invocation, not an asserted improvement in a final tower height. Other losses may dominate. This question is valuable for the repository even if the final asymptotic class does not change, because it would identify which local constants are actually bottlenecks.

Research question V.152.7 (Beyond the homogeneous positivity class). Can vertex-dependent kernels or adaptive selection beat the $\alpha^s \eta^{s-1}$ certificate scale, while retaining a usable exceptional-set estimate?

Theorem II.172.2 makes this a precise boundary question. A candidate improvement should state exactly which hypothesis of that theorem it leaves, and should exhibit either a new positivity mechanism or information about the configuration family absent from the scalar first-moment bound.

Research question V.152.8 (Formal proof and executable certification). Can the one-feature relation theorem, exact positive-kernel moment, and score selection theorem be formalized as reusable components, with the finite checks generated from the same definitions?

A practical first milestone is the algebraic universal-relation theorem, which does not require probability or analysis. A second is a finite Fourier moment theorem parameterized by a nonnegative coefficient family. A final milestone is the explicit arithmetic corollary with all characteristic and size inequalities present in the formal statement.

V.152.1 Integration recommendation

Place this package in a new subdirectory named `fejer-arrangement-selection` of the existing `local-quantitative-refinements` collection under `Research/GowersSzemerédi`. The accompanying README provides an index entry and a proof-status summary. Keep the original paper’s transcription unchanged: the new theorem should be a separate, explicitly cited refinement. Later Lean integration and any downstream constant propagation should be separate changes with their own verification record.

V.153 Conclusion

Source 40, delivered Section 12.

The repeated degree-one Riesz-product selection in Section 15 can be replaced by one high-bandwidth Fejér selection. The main quantitative benefit is a reduction from the exponent $s2^{s-1}$ to powers s and $s - 1$ in the input abundance and error parameters, together with a smaller explicit field-size requirement in the eight-arrangement case. The algebraic input is also compressed: the full coordinate product alone has precisely the intended universal relation, and all other bounded relations are controlled by a polynomial zero-count.

The argument is local, fully stated, and separates exact identities from conservative estimates. Its principal technical safeguards are the explicit treatment of repeated vertices, the characteristic condition preventing coefficient aliasing, and the accounting of exceptional bad arrangements before extracting a subset. The positive-Fourier obstruction identifies a genuine limit of this particular method, while the torsion formula and the research agenda identify concrete directions for further work.

V.154 The pinned source and classical context

Source 41, delivered Section 8.

V.154.1 The source implication in every degree

The precise target in the repository [193] is `corollary_5_8` in `Section05.lean`, lines 242–259, at commit

`c94998a5db18ba63c815a3bf6c3877a0fde91930`.

The source uses $K_k = (k!)^2 2^{(k+1)^2}$ and asks for length at least $S^{1/K_k}/8$ with gain $\alpha/16$. Its `PolynomialOn k Finset.univ` means a coefficient sum indexed by `Fin (k+1)`, exactly the notion used here. Its `ModAP.IsProper` says carrier cardinality equals parameter length, equivalently injectivity of the finite parameter map. It permits wrapping without repetition. Its density identity identifies the balanced function with $\mathbf{1}_A - \delta$.

Empty parents are syntactically allowed by the source. If one is empty, two-comparability forces every parent empty, and coverage makes A empty, contradicting positive correlation. The source’s remaining modulus $N = 1$ is also vacuous, since the balanced function vanishes. Thus our nonempty-parent and $N \geq 2$ presentation adds no nonvacuous hypothesis.

Since $D_1 = 4 \leq 16 = K_1$, and $D_k = K_k/4 + 1 \leq K_k$ for $k \geq 2$, $S \geq 1$ implies that (F.46.4) dominates the required length in every degree. The gain $\alpha/3$ dominates $\alpha/16$. There is no concluding parent-containment requirement in the source.

The target is a `Prop`-valued catalogue declaration. The argument proves its mathematical interface in writing; it does not add a theorem to the repository, run its build, or change its formalization status. Source links and fingerprints are recorded in `SOURCES.md`.

[write] 7 October 2026; sources 40–46. Checked at HEAD 6d4942222. Section05.lean and Definitions.lean have the blobs 882b17a and 97113b7 recorded in SOURCES.md (shipped as 41-all-degree-increments-SOURCES.md); corollary_5_8 occupies lines 242–256 (the cited range 242–259 runs into the docstring of lemma_5_9; immaterial); it now has the exact companion corollary_5_8_holds (Proofs05ThresholdFreeLocalization.lean, line 86). The pinned statement is described accurately above (gain $\alpha/16$, length $(N/M)^{1/K_k}/8$, Polynomial0n k as a coefficient sum, ModAP.IsProper as injectivity).

V.154.2 Primary literature and what is being claimed

Gowers [1], Section 5, develops the classical chain from Weyl recurrence to polynomial localization and local density increments. Its Corollary 5.8 on printed page 493 displays gain $\alpha/8$, while its displayed proof ends with $\alpha/16$. The pinned edited source uses $\alpha/16$. The present argument supplies explicit recurrence and localization budgets and uses a subgroup-coset variance alternative to handle the low-correlation branch. It does not import an internal large-threshold statement without checking that the available parent length pays for it.

Maynard [72], Theorem 1.1 and Corollary 1.2, gives coefficient-uniform simultaneous recurrence for systems of real polynomials vanishing at zero. For m polynomials of degree at most d , the corollary has fractional parts bounded by $\ll_{d,m} X^{-c_d/m}$ for some $1 \leq n < X$. The introduction explicitly leaves degree and dimension dependence of implied constants uncontrolled. That result gives important context, but its implicit constants are not substituted for the explicit constants proved here. In the cited arXiv version the Weyl estimate is Lemma 6.1; numbering need not agree with later versions.

The repeated-differencing and recurrence-to-localization mechanisms are classical. The explicit constants here are deliberately weak, and no literature-priority or optimality claim is made for either those constants or this local-cover formulation. This local theorem alone neither repairs other source sections nor produces an ordinary-integer rectification theorem or a better global quantitative Ramsey or Szemerédi bound. Those require separate statements and compatible downstream arguments.

V.155 Reproducibility and further questions

Source 41, delivered Section 9.

V.155.1 The role of exact diagnostics

The companion uses only the Python standard library, integer arithmetic, and exact rational arithmetic. Its bounded diagnostics check the recurrence and partition budget identities, ordered-divisor combinatorics, differencing coefficients, subgroup-coset second moments, removal of nonproper directions, residue-chain chunking and resplitting, modular polynomial pullbacks, cover mass, and rounding regressions. The JSON output reports the actual finite ranges and limitations. Small rational recurrence examples test an interface; they do not prove recurrence for arbitrary real coefficients.

The infinite-degree induction and the real-parameter recurrence are proved in this article. The program does not enumerate astronomical theoretical partitions or build a theorem-scale integer merely to display its digits. External numeric inputs have explicit runtime caps, active under optimized Python. Tests and diagnostics run both normally and with `-O`, and deterministic diagnostic stdout must agree.

The offline builder verifies an exact source inventory and, for a packaged distribution, manifest hashes. It runs bounded checks, rebuilds the PDF using an installed trusted TeX toolchain with shell escape disabled, and creates a deterministic ZIP with fixed order, timestamps, and permissions. It rejects unexpected source members, aliases, unsafe output locations, and oversized subprocess output. These are reproducibility and regression controls for trusted source, not a sandbox for hostile TeX or Python, and not a guarantee against an actively malicious concurrent filesystem owner. Hashes show byte agreement, not authorship or mathematical truth.

The package contains only the article, PDF, original exact diagnostics and tests, the builder, public documentation, and the manifest. Third-party papers are linked rather than bundled, and private review material is excluded. The ZIP digest is supplied outside the archive. Reproduction commands, environmental requirements, caps, and trust boundaries appear in `README.md`.

V.155.2 Focused further questions

1. **Improve the recurrence with fully explicit constants.** Can modern monomial or simultaneous recurrence estimates be made effective at all relevant small scales and then inserted into the exact recursion? An improved asymptotic exponent alone does not identify the required prefactor or cutoff. The conditional witness in Remark I.64.4 is not such an improvement.
2. **Reduce localization cost.** The induction suppresses the leading coefficient, then applies the lower-degree result at target $2L$ and resplits. Can one reduce the exponent e_k or constant c_k , perhaps by optimizing the allocation of phase error, while retaining every point and keeping an integer-scale proof?
3. **Optimize the final selection without losing scale control.** The exact-budget corollary exposes what the present recursion actually proves. Can a sharper variance estimate or a direct correlation argument improve its floor and prefactor, or trade the uniform gain $\alpha/3$ against a longer progression? Any comparison should retain singleton and boundary cases.
4. **Strengthen geometric conclusions.** When can the low-correlation branch be made parent-contained? When can a modular output be rectified to an ordinary integer progression with a useful length? Neither property follows from properness alone over arbitrary cyclic moduli.
5. **Formalize and connect the interface carefully.** The finite-group mass and variance lemmas, residue partitions, explicit recurrence, and numerical induction could be formalized as separate components and composed into the pinned predicate. Beyond that local task, what downstream statements would be needed before claiming a global quantitative consequence? This report provides no kernel-checked implementation or global conclusion.

V.156 Relation to prior arguments and further questions

Source 42, delivered Section 8.

V.156.1 What is strengthened and what is classical

The inspected variance argument in Report277 averages over all cyclic steps. Off-diagonal correlations are controlled by subgroup-coset averages; the nonproper directions are then removed using the estimate

$$\eta \leq \frac{L(L-1)}{2N}.$$

At $N \geq L^3$, this yields the gain $(1-\delta)/(2L)$. The present proof changes the step distribution before taking moments. Its only nonproper sampled direction is zero, with exactly controlled mass. That change leads to the quadratic condition, the arithmetic refinement, and the equality examples given here.

Autocorrelation positivity, finite Fourier inversion and Poisson summation, variance-to-range conversion, and positive-definite packing estimates are classical. The role of Kolountzakis–Révész is particularly direct: Proposition I.71.1 is an instance of their Theorem 2, rather than a new packing inequality. Bhatia–Davis supplies the standard name for Lemma I.69.2, whose elementary proof was included.

Fox, Xu and Zhou [183] study discrepancy of modular arithmetic progressions. Their Section 4, especially Lemmas 4.2–4.3 and Corollary 4.4, is a close neighboring second-moment mechanism: it uses Fourier weights, subgroup sums and the behavior of short-order directions. Their objective is two-sided discrepancy over

varying progression lengths. The fixed-length one-sided density statement and the explicit prime-quotient equality calculation are the claims isolated in this report. The comparison is a bounded source screen, not a proof of priority.

Lewko [101, Lemma 2 and Section 5] gives a close contemporary use of autocorrelation sums of squares, progression second moments and explicit zero-step removal in a prime cyclic group, with an additive-energy hypothesis. These are useful neighboring mechanisms; their use does not by itself supply the unconditional all-modulus coefficient proved here. In the pinned ProveIt research article [194], the source-10 contribution already records the elementary implication $M \geq \delta + \text{Var}(X)/\delta$. The present report does not claim that conversion as a new argument.

Gowers [1, Section 5] provides the polynomial localization and transfer context used in the earlier reports. The present standalone variance improvement does not reprove or improve their all-degree phase budget. A global quantitative Szemerédi or Ramsey improvement would require a separate, complete iteration or transfer argument. Likewise, the present written proof is not a Lean or other kernel-checked certificate.

V.156.2 Focused open directions

The theorem leaves several concrete questions.

1. *Fixed-density sharpness.* For δ bounded away from zero and one, what is the best uniform coefficient under $N \geq CL^2$, including $C = 1$? Theorem I.73.1 supplies a positive endpoint coefficient but does not prove it optimal. The sparse equality family does not decide the fixed-density minimax problem.
2. *Arithmetic pairs beyond the equality family.* For fixed N, L and permitted cardinality $|A|$, when is the arithmetic bound attained? The support optimum and the full extremal set problem are distinct. Nonzero off-diagonal squares and strict range inequalities may force additional gain.
3. *General finite abelian groups.* Theorem I.69.3 reduces one route to finding a large support with no nonzero short-order difference. A short-order subgroup supplies a packing upper bound, but there need not be a cyclic interval attaining it. Determining the best support and positive-definite weight parameters in more general groups is a separate problem.
4. *Containment and localization.* Can comparable arithmetic gains be obtained while keeping the output inside a specified parent or a cover? The global translation average used here does not supply that constraint. Any nonwrapping integer conclusion likewise requires an additional argument.

These questions concern added structure or stronger conclusions. None is required for the proofs above.

V.157 Reproducible bounded diagnostics

Source 42, delivered Section 9.

The archive contains the full TeX source and PDF, a standard-library Python companion, tests, a deterministic builder, a source note and a SHA-256 manifest. Exact integer and rational diagnostics check proper-window maxima, covariance squares, variance decomposition, admissible supports, bounded functions, prime-quotient examples, endpoint sliding and coset identities, and integer rounding. The full profile includes all 81,792 nontrivial triples (N, L, A) with $2 \leq N \leq 12$ and $2 \leq L \leq N$.

These are bounded diagnostics, not universal proofs, extremizer classifications or formal certificates. In particular, scalar affine examples do not certify an actual phase-correlation premise. The unrestricted mathematics is justified by the written proofs above. Public numerical helpers have explicit computational caps and runtime validation that remains active under Python optimization.

The builder runs the default diagnostics and both test suites normally and with optimization, requiring matching standard output. It compiles with shell escape disabled, uses fresh external output directories and fixed archive metadata, checks unchanged source bytes, and verifies the public inventory and hashes. A

pinned distribution must reproduce its PDF bytes exactly. The README gives all commands, caps and the fuller trust boundary. The builder assumes trusted sources and installed tools; it is not a sandbox for hostile code. Matching hashes certify bytes, not mathematical correctness or priority.

V.158 Relation to Gowers’s argument and the formal project

Source 43, delivered Section 10.

V.158.1 A precise crosswalk, not a replacement of hypotheses

The relevant formal definitions at the comparison snapshot are in `Sections17_18.lean`. The definition `proposition_17_2` has ambient group $\mathbb{Z}/N\mathbb{Z}$ with N prime, a disc-valued function, a frequency predicate `IsMultilinear`, and an explicit unit condition on $(k + 1)!$. It concludes the existence of a degree-at-most- $k + 1$ phase preserving the stated lower bound on the sum of squared zero-frequency derivative coefficients after twisting.

There is an important naming distinction. At the pinned snapshot, `Definitions.lean` defines `IsMultilinear` as a sum of square-free monomials *including constant and lower-order terms*. It is a multiaffine predicate. Our T is genuinely homogeneous multilinear in all its slots. A general multiaffine selector is not covered merely by taking its top homogeneous component.

For the homogeneous cyclic high-characteristic specialization, the usual diagonal primitive is available because the factorial is invertible. The present theorem addresses an obstruction that appears on vector spaces in characteristic two. It certifies the additional identities needed for a nonclassical primitive and preserves energy exactly once they hold. This is a local analogue in a broader class of ambient groups, with its own high-energy hypothesis, not a theorem that uniformly dominates Proposition 17.2 for every input covered by the repository.

The displayed degree correction from k to $k + 1$ in the source paper is already recorded by the repository. It is not a correction originating in this manuscript. Likewise we do not alter the local partition constants in Proposition 17.7 or any subsequent density iteration.

Target	Input retained or changed	Output here
Gowers global phase removal	Homogeneous selector on $\mathbb{F}_2^n$; new high-energy condition	Exact nonclassical phase removal, Theorem IV.38.4
Prior odd-characteristic symmetry criterion	All finite abelian groups; arithmetic radical profile	Theorem IV.36.4, including the sharp Boolean threshold
Tidor’s bilinear estimate	Same three-function finite-field problem; also finite abelian groups	Optimal fourth power instead of eighth, Theorem IV.40.1
Tidor’s symmetric trilinear repair	Same seven-function hypothesis	Codimension at most $\lfloor 2\log_2(1/\delta) \rfloor$, Theorem IV.40.3
One-function cubic repair	Prescribed symmetric cubic selector	Exact energy and optimal codimension, Theorem IV.39.4

V.158.2 An executable algebraic route

Given the coefficient table of a d -linear Boolean form, the certification and integration steps can be implemented exactly. First check symmetry on coefficient tuples, using adjacent transpositions. Once symmetry is known, the repeated-variable defect is multilinear, so checking it on basis arguments suffices. The test uses the coefficients

$$T(e_i, e_i, e_j, e_{a_1}, \dots, e_{a_{d-3}}) + T(e_i, e_j, e_j, e_{a_1}, \dots, e_{a_{d-3}}).$$


If they vanish, read one coefficient per nonempty support and construct (IV.37.2). A dense input has n^d coefficients; checking adjacent symmetries costs $O(dn^d)$ field operations, and the repeated-variable tests cost $O(n^{d-1})$ coefficient operations after symmetry has been established. The construction outputs $\sum_{s \leq d} \binom{n}{s}$ rational monomials, with denominator at most 2^d . These are algebraic operation counts, not a claim of sublinear access to a dense tensor.

For a cubic, form the $n \times n$ matrix $\Omega_{ij} = T(e_i, e_i, e_j) + T(e_i, e_j, e_j)$. Symplectic Gaussian elimination gives its rank, an optimal isotropic subspace, and the linear forms in (IV.39.6). This costs $O(n^3)$ field operations with a straightforward implementation. The normal form then provides a global phase extremizer and an optimal restriction primitive.

An observed energy lower bound is a different kind of input. If function values are numerical complex numbers, an interval or exact cyclotomic certificate is needed before a strict threshold can be invoked. An unchecked floating-point comparison is not a proof of (IV.38.8).

V.158.3 Proposed Lean boundaries

No new theorem in this package has been checked by Lean. A useful formalization order is the following mathematical dependency chain:



A first implementation can work solely over $\mathbb{F}_2^n$ and use matrix symplectic elimination. The full finite-abelian Lagrangian theorem can be added independently. The polynomial value group must include $2^{-d}\mathbb{Z}/\mathbb{Z}$; using only ZMod 2 would encode the wrong integration problem.

The relevant finite-dimensional analytic obligation is explicit: a positive operator commuting with the projective translations has nonzero eigenspaces of dimension at least d_B . Once its trace is computed, the sharp constant follows. There is no appeal to an unformalized inverse theorem hidden in that step. The quartic proof additionally needs the finite monomial identity and a disc-boundary maximum argument, or an equivalent real-algebraic certificate. The file `FORMALIZATION.md` expands these proposed interfaces without supplying placeholder Lean proofs.

V.159 Exact checks, provenance, and limitations

Source 43, delivered Section 11.

V.159.1 What the supplied computation establishes

The companion `checks/verify.py` uses only Python's standard library. All mathematical comparisons use integers, rational numbers, or exact arithmetic in $\mathbb{Q}(\sqrt{2})$. The fixed pseudorandom seed is 20261006; the randomized tests are exact evaluations of selected cases, not probability-based mathematical proofs.

Check	Finite scope	Result
Symmetric cubic census	16 tensors; 512 phase functions each	8,192 exact energies obey and attain $2^{-\text{rank } \Omega/2}$
Nonsymmetric cubic census	240 tensors; 64 phases each	15,360 exact energies; maximum 3/4
Primitive formula	All support coefficients on $\mathbb{F}_2^2$, degrees 2–5	217,600 pointwise top- and next-derivative checks
Pointwise shear	Gaussian rational values, including zero and nonunit amplitudes	15,360 exact pointwise equalities
Composite orthogonality	93 cyclic-pair coefficient cases	Radical sizes and equality factors agree exactly
Mixed estimates	Explicit sharp triangle and 1,024 sampled seven-function cases	Fourth-power equality and rank-bound checks
Quartic identity and phases	All 1,024 cube terms; 4,096 phase functions	29-monomial certificate and exact phase-grid maximum 11/16

The cubic phase census fixes $f(0) = 1$ and uses eighth roots of unity on the other three points. The nonsymmetric census uses fourth roots. The quartic phase census uses sixteenth roots. The monomial identity is checked before substituting phases, so it is not restricted to a phase grid. Additional tests verify the endpoint constructions, canonical cubic blocks through rank six, and all subspaces of $\mathbb{F}_2^n$ for $1 \leq n \leq 4$ in the isotropic repair calculation.

These checks detect sign, normalization, characteristic, and endpoint errors. They do not replace the proofs for arbitrary dimension, arbitrary complex amplitudes, or all finite abelian groups. In particular, the proof of the continuous quartic maximum is the subharmonic and trigonometric argument, not an inference from the 4,096 phase values.

V.159.2 Source audit and novelty boundary

The file `data/source_audit.json` records the exact repository revision and blob identifiers used for the source comparison. The prior manuscript [197] was read from the user’s Library; its source labels, rather than unstable printed page numbers, identify the questions addressed. The repository’s aggregate local-refinement README was also checked to avoid repeating the density, carry-pattern, cube-count, and Bohr-radius projects already represented there.

There are three distinct statuses. The mathematical theorems in this manuscript have supplied proofs. Their selected finite instances have executable exact checks. They have not been independently refereed or proved in a proof assistant, and the targeted literature search does not establish publication priority. The package’s suggested repository placement conveys no additional formal status.

V.159.3 Limits that affect downstream use

The universal Boolean threshold is a high-energy criterion for an already supplied multilinear selector. It is not an inverse theorem that constructs a selector from an arbitrary positive Gowers norm. The exact cubic law remains useful at all positive energies, but still assumes an exactly symmetric global cubic tensor. Producing or repairing that symmetry may itself be a substantial part of an application.

Uniform translation averages are essential to the projective twirl. Restricting the averaging set to a progression or a Bohr set generally breaks the commutation argument. No boundary-free local analogue is claimed. Similarly, multiaffine lower-order frequency terms cannot simply be discarded, and the higher-dimensional quartic maximum is not settled by its two-dimensional computation.

V.160 Further research questions and topics

Source 43, delivered Section 12.

The first questions below are sharply specified continuations of the proved results. The later ones concern their use in a complete inverse or density-increment argument. An assertion marked as a question is not used elsewhere as a theorem.

Research question V.160.1 (Does the quartic extremum stabilize?). Is $\mathcal{M}(C_{4,n}) = 11/16$ for every $n \geq 2$? Equivalently, by Corollary IV.41.4, is $11/16$ the exact universal energy threshold for symmetric nonintegrable quartics? A counterexample must use the extra coordinates in the optimizing function, since the tensor itself has already been reduced to two coordinates.

[write] **7 October 2026; sources 47–53.** *Answered by source 47* (Section IV.42): $\mathcal{M}(C_{4,n}) = 11/16$ for every $n \geq 2$, and $11/16$ is the exact threshold for symmetric nonintegrable quartics in every dimension (its Theorems F.52.1 and F.52.2; in every degree $d \geq 3$ the canonical maximum is $1 - (d+1)/2^d$). Source 47 cites this source by its library file name.

A useful first target is $n = 3$, retaining the exact shear identities rather than treating the sliced functions as independent optimizers. The interval $11/16 \leq M_{4,n} \leq 3/4$ is proved, not conjectural; constancy at the lower endpoint is a proposed conjecture here.

Research question V.160.2 (Exact mixed cubic repair). For a prescribed symmetric cubic T , determine $\sup_{|b_i| \leq 1} \Gamma_T(b_1, \dots, b_7)$. Does it depend only on $\text{rank } \Omega_T$, as the one-function extremum does? Is the upper bound $2^{-\text{rank } \Omega_T/4}$ attainable, and is coefficient two in (IV.40.8) optimal?

Sharpness of the triangle estimate does not settle simultaneous attainment by all the sheared seven-function slices.

Research question V.160.3 (Near-extremizers beyond their amplitudes). If $\mathcal{E}_T(f) \geq (1 - \varepsilon)2^{-m}$, classify f up to a cubic integrable twist and the projective eigenspace structure. Can a dimension-independent distance bound be proved for an appropriate explicit family of extremizers? The amplitude estimate (IV.39.13) is available; phase stability is not.

Research question V.160.4 (Higher-degree defect geometry). For $d \geq 5$, characterize the multilinear families $\mathbf{z} \mapsto \Omega_{\mathbf{z}}$ that actually arise from symmetric d -linear tensors. Which compatibility identities generalize (IV.41.6)? Can low-dimensional defect images reduce the sharp energy problem to finitely many canonical tensors?

The existing support bound treats an arbitrary multilinear family; using the additional compatibility is the natural place to improve it.

Research question V.160.5 (Odd small characteristic). For $p \leq d$, find a spectral identity detecting the p -repetition transfer conditions of nonclassical integration. Determine whether $1 - (1 - 1/p)^{d-1}$, already the sharp symmetry threshold, also certifies all the extra integrability identities. The Boolean shear exploits $2a = 0$ in an exact way; a p -cycle substitute is needed.

This asks for certification of known algebraic integration conditions, not for rediscovery of nonclassical polynomial theory.

Research question V.160.6 (Composite-group primitives). On a general finite abelian group, characterize which symmetric multiadditive $T : G^d \rightarrow \mathbb{T}$ have a degree-at-most- d primitive, and derive sharp energy certification using elementary-divisor data. Can the primary factorization of the symmetry profile be lifted to an equally explicit integration theorem?

The present paper settles the finite-abelian orthogonality and symmetry parts, not this full integration problem.

Research question V.160.7 (A boundary-sensitive local twirl). Replace the global translation average by an arithmetic box or a regular Bohr set. Can its failure to commute with projective translations be bounded in a norm that preserves a useful rank factor? A successful statement should give explicit boundary and normalization errors compatible with Gowers’s localized phase step.

An estimate that simply repeats the global proof on a restricted set would miss its crucial invariance hypothesis.

Research question V.160.8 (Obtaining the high-energy selector). Identify a nontrivial inverse-theorem regime in which a global multilinear selector can be produced above the certification threshold, or in which the exact cubic repair law is useful below it. Track the selector, symmetrization, phase removal, and final density increment together before claiming a better global bound.

Research question V.160.9 (Multiaffine frequency offsets). Can lower-order components of a multiaffine selector be removed with an optimal quantified loss, or absorbed into a larger exact model? A satisfactory theorem must account for the constant and lower-order terms present in the repository’s `IsMultilinear` predicate.

Research question V.160.10 (Formal certificates for extremal identities). Formalize the projective eigenvalue-multiplicity argument and the Boolean primitive formula, then verify the cubic exact-rank theorem. For the quartic result, produce a small proof object combining the 29-monomial identity with the continuous optimization bound, rather than treating the finite phase census as a certificate of the maximum.

V.161 Conclusion

Source 43, delivered Section 13.

The additional characteristic-two integrability obstruction is not merely detectable near perfect energy. For symmetric cubics, its alternating rank determines the entire extremal value and the exact cost of removing the obstruction by restriction. For arbitrary degree, a division-free spectral argument plus a Boolean shear forces the full set of nonclassical integration identities at a sharp universal threshold, after which phase removal is lossless.

The same orthogonality mechanism is useful outside the one-function setting: it gives an optimal fourth-power bilinear estimate and a smaller codimension bound under the published seven-function repair hypotheses. The quartic analysis provides a second structural result: an exact two-dimensional value and a reduction of the general optimization to one explicit tensor. The remaining ambient-dimension question is isolated, not hidden in an assumed tensorization step.

These results are suitable as independent research additions beside the formal Gowers development. Their integration into a global Ramsey or Szemerédi estimate requires the further selector and localization work described above.

V.162 Verification, source comparison, and formalization targets

Source 44, delivered Section 11.

V.162.1 What was actually executed

The supplied verifier uses Python’s standard library and exact integers and rational numbers. Its default run computes through $d = 10$. It also performs checks which do not use the fixed-grid representation. The saved record is `data/verification.json`; every asserted check passed in the packaged run.

First, all 512 grid subsets are considered and all 458 spanning supports are reconstructed from their affine-value assignments. The universal eight-shape coefficient identity is checked over the complete support

lattice, grouped by exact rational affine equivalence. This is stronger than comparing the compact formula with a few numerical values. The walk recurrence, exact-support Möbius transform, flat multiplicity formula, and both pair-consistency identities are then checked for $d = 1, \dots, 10$.

Second, the verifier generates every integer-balanced ternary normal for $d = 2, 3, 4$, canonicalizes its rational restriction, and groups every pair by primitive exterior-product coordinates. These independent pair-span histograms match the grid answer. Modular two-row reduced echelon forms at $p = 17, 53$ give the same histograms. These calculations check examples of the determinant theorem; they do not replace its all-dimensional proof.

Third, affine ternary Boolean maps are exhaustively checked for $k = 1, \dots, 8$ at $p = 5, 7, 11$. The number of maps is $3 + 6k + 4\binom{k}{2}$, and none has more than two active slopes. The exact signed-collision counts are checked directly for $m = 4$, $p = 5, 7$, and $k = 1, 2, 3$. In characteristics two and three the quotient counts and the regular cross-section formula are checked using the prime-field vector-space models of $\mathbb{F}_4, \mathbb{F}_8, \mathbb{F}_9$. Only addition and prime-field scalar multiplication are needed for these checks; no assertion about an unimplemented field multiplication routine is made.

Finally, the seven-step walk recurrence is evaluated through $d = 64$. The scaled energy is compared numerically with the proved leading constant and its first correction at $d = 8, 16, 32, 64$. These floating-point ratios are explicitly labelled diagnostics. All preceding counts and certificate tests use exact arithmetic.

The run does *not* enumerate every full arrangement, verify a new Szemerédi bound, or execute a Lean build. A finite check is not called an independent mathematical review. The universal statements rest on the proofs; the eight-shape compression has the separately delimited finite rational component described above.

V.162.2 Contribution ledger

Statement	Relation to inspected baseline	Evidence in this package
First-order hyperplanes and H_d	Earlier manuscripts [199, 200]	Counting interpretation recalled here.
Direction inventory and non-core cubic bound	Earlier manuscript [200]	Reproved in Section II.177.
Width classification for every $p \geq 5$	Removes the earlier $p > 2k + 2$ cutoff	Direct translation-growth proof; finite checks.
Exact core products	Strengthens the earlier pairwise core bounds	Exact sequential counts in Sections II.176 and II.182.
Uniform rank-two characteristic cutoff	Makes the relevant incidence coefficient stable independently of k	Integer determinant proof.
Fixed-grid formula and polynomial-operation algorithm	Answers the explicit b_2 computation question	Weighted double count and walk recurrences, independent of compact certificate.
Eight-shape compression	Further simplification of the new formula	Complete finite rational coefficient certificate.
$d = 8$ second coefficient	Evaluates the previously unspecified invariant	Exact recurrence output, consistency identities, and proved reduction.
Deficit asymptotic	New secondary scale for the invariant	Seven-point support classification and Gaussian integral expansion.

V.162.3 Repository integration

A separate research directory is preferable to silently changing the statement of Lemma 15.4:

`Combinatorics/Ramsey/Research/GowersSzemerédi/finite-grid-core-avoidance/`

The intended interface is `GeneralArrangement.IsDegenerate`, `arrangementMoment`, and `degenerateGeneralArrangementCount` in the inspected `Sections14_15.lean`. The article's k, d agree with cube dimension and arrangement multiplicity there; $m = 2d$ counts cross-sections, and $v = m2^k$ counts labelled vertices. The zero-side convention and the modular-multiple exclusion are preserved.

For a first formalization, the characteristic-uniform width lemma and the signed-collision product are the smallest independent targets. Both can be expressed with finite sets and elementary linear maps. Next isolate the rank-two determinant comparison and the weighted profile-pair double count. A verified walk recurrence and a finite rational-affine checker can then certify the numerical specialization. The full geometric language of characteristic polynomials is not required to prove the finite-field second-order bound.

No placeholder Lean code is included. The package does not modify the repository or its formalization ledger. `FORMALIZATION.md` records proposed theorem interfaces and dependencies, not completed proof-assistant results.

V.163 Further research questions and proposed next targets

Source 44, delivered Section 12.

The first two questions are the most direct mathematical continuations. The last one is the route toward a genuine improvement of a global Ramsey bound. These are proposed research questions of this manuscript; no claim is made that every formulation has previously appeared as a named open problem.

Question V.163.1 (Third Whitney coefficients from a fixed grid). *Can the third Whitney coefficient be expressed by balanced supports in $\{-1, 0, 1\}^3$, with weights determined by a finite rank-three incidence type, and evaluated in polynomially many arithmetic operations in d ?*

The rank-two proof uses ordered bases of a normal span and divides by the number of such bases through a flat. In rank three, the Möbius coefficient depends on the full rank-two incidence structure inside the span, not only on its number of hyperplanes. The finite grid has 27 points. A useful answer should specify the correct weight, exploit symmetry before enumerating its supports, and distinguish arithmetic complexity in d from the large universal preprocessing constant.

Question V.163.2 (The third term of the actual degeneracy count). *Which non-core directions contribute at order q^{-3} on the regular region, and what are the intersections of their irreducible components?*

A third Whitney coefficient alone is not enough to answer this. The rank argument separates directions with one extra cross-section equation from directions forcing two independent rank-two base systems. The latter cost q^{-4} . A structural classification of the former could replace the entire residual inventory by a smaller third-order family. The exact core product already settles all core intersections, so it need not be reanalyzed at that stage.

Question V.163.3 (Exact recurrences for the eight energies). *Find short polynomial-coefficient recurrences, with proofs, for the energies in (II.180.8), and determine whether their linear combination B_d has a substantially simpler minimal recurrence.*

The fixed-step walk model supplies exact data without guessing. A recurrence must be certified by a coefficient identity or a valid summation certificate, not just fitted to an initial segment. Such a result could reduce both storage and the number of arithmetic operations, and could supply all-orders asymptotics by an independent route.

Question V.163.4 (Sharp exceptional characteristics). *For each d , which primes actually change $b_{2,d,p}$, and is the universal bound 16 in Theorem II.178.1 optimal after the balancing relations are imposed?*

The determinant estimate intentionally ignores much of the structure. A prime can divide a possible minor without being realized by a balanced triple in the actual hyperplane arrangement. Intermediate characteristics also introduce additional balanced profiles when $p \leq 2d$; those are a separate effect from collapse of rational incidences. A complete answer should distinguish these mechanisms.

Question V.163.5 (Sharpness of the universal dimension barriers). *How large can k be for a nondegenerate d -arrangement to exist over a fixed field? How close are the necessary conditions in Corollary II.176.4 to sufficient ones?*

The core laws give exact capacities for the paired summaries. Realizing summaries within those capacities does not rule out non-core relations. A construction with certified absence of every ternary moment relation, or a stronger obstruction from the next coefficient layers, would settle an existence problem beyond the classified core.

Question V.163.6 (Relative and nonuniform arrangement distributions). *Can an analogue of the core product control arrangements sampled from a dense set, rather than the full ambient parameter space?*

The present product relies on uniform independent paired summaries. A relative version needs explicit hypotheses on their conditional concentration and collision probabilities. Total variation estimates, weighted sequential exposure, or a mixing condition on the base fibres are possible starting points. An ambient bound multiplied by the inverse collection density is not the desired improvement.

Question V.163.7 (Larger coefficient alphabets). *For $\{-M, \dots, M\}$ in place of $\{-1, 0, 1\}$, can one obtain a fixed-grid formula for the second incidence coefficient and an exact classification of the low-degree core events?*

The weighted profile-pair argument suggests a grid with $(2M + 1)^2$ points, but the number of alphabet representatives of an affine class is no longer determined solely by whether two or three values occur. The core event may involve more than signed collisions. A satisfactory extension should keep those multiplicities and the prime-dependent translation symmetries explicit.

Question V.163.8 (A small independently verified certificate). *What is the smallest trusted checker needed to certify the eight-shape formula and the $d = 8$ coefficient inside the repository's formal system?*

The general double-counting theorem should be proved once. The numerical part can then be limited to integer walk arrays, rational affine-map checks on at most nine points, and finite sums of rational weights. This suggests a certificate substantially smaller than enumerating the millions of hyperplanes in the $d = 8$ arrangement.

Question V.163.9 (Complete Section 15 propagation). *After inserting the exact principal term and the improved incidence coefficient into the corrected selection argument, which parameter actually controls the next inverse step and the final density increment?*

This requires a complete ledger of retained arrangement mass, the non-core error, repetition counts in the selection kernel, field-size thresholds, and progression widths. The negative second coefficient can improve a local estimate without changing the final asymptotic bound. Only a full calculation through the later stages can establish a global quantitative improvement.

V.164 Conclusion

Source 44, delivered Section 13.

The computational obstruction in the second degeneracy coefficient is not the number of ambient hyperplanes. Two ternary profiles can be recorded by their coordinate pairs in a fixed nine-point grid. Exact representative weights turn this observation into a double-counting identity and a polynomial-operation algorithm. The previously unspecified coefficient in the 8-arrangement setting is thereby evaluated, with a uniform characteristic range and a checkable finite certificate for a compact formula.

The core geometry also admits an exact count. In characteristics at least five it is a problem of avoiding a finite offset set and repeated sign orbits. In characteristics two and three it is a linear-independence problem

in a quotient vector space. These descriptions eliminate the intermediate-characteristic gap and give exact finite products, universal dimension obstructions, and distinct growing-dimension laws.

The boundary of the advance is explicit. Non-core directions still contribute a cubic remainder, the eight-shape compression includes a finite computer-assisted component, and no global Szemerédi estimate has been propagated from the local refinements. The new formulae make those remaining tasks more focused rather than disguising them as completed results.

V.165 Reproduction and package contents

Source 44, delivered Appendix A.

The archive contains `article.tex`, its compiled `article.pdf`, a `Makefile`, `README.md`, `SOURCE_AUDIT.md`, `FORMALIZATION.md`, the standard-library verifier, and exact data and certificates. From the package directory,

```
python3 code/verify.py --max-d 10
latexmk -pdf -interaction=nonstopmode -halt-on-error article.tex
```

reconstructs the checks and the PDF. The verification program requires Python 3.10 or later and must not be run with assertion-disabling optimization. It performs no network requests and makes no repository changes. The TeX source uses ordinary packages and an embedded bibliography; it requires no shell escape, external figures, or external bibliography service.

The machine-readable verification record distinguishes exact checks from the floating-point asymptotic diagnostics. The rational normal lists for $d = 2, 3, 4$ are included to make the independent pair-span check inspectable. No fonts, downloaded papers, or earlier research manuscripts are redistributed in this package.

V.166 Further research questions and a proposed programme

Source 45, delivered Section 6.

The following questions isolate the remaining mathematical gaps. They are proposed directions, not consequences of the preceding theorems. The first group concerns finite classifications, the second concerns the best stability ranges, and the third concerns the functional extremizers that control cube counts.

V.166.1 The next energy levels

Question V.166.1 (All third-level extremizers). *Classify every m -element torsion-free set satisfying $E(A) = M_m - 8(m - 3)$. For sufficiently large m , are the only models, up to translation, dilation, and reflection,*

$$\{0, 1, \dots, m\} \setminus \{2\} \quad \text{and} \quad \{0, 2, 3, \dots, m - 1, m + 1\}?$$

What is the smallest valid threshold, and what additional families occur below it?

[write] **7 October 2026; sources 54–59.** *Answered by source 57* (Section II.291): for $m \geq 7$ the only models are the two displayed, up to affine maps, and 7 is the optimal threshold, the one further class at $m = 6$ being $\{0, 1, 2, 4, 5, 6\}$ (its Theorem II.292.7; item II.53).

[write] **7 October 2026; sources 60–69.** *Also answered by source 66* (Section II.296), with the threshold $m \geq 8$ for its three oriented templates (source 57’s two affine classes), and the fourth level $M_m - 8m + 20$ for $m \geq 11$ (its Theorems II.307.1–II.308.2); source 57’s threshold $m \geq 7$ is the better one.

The second displayed model also has defect $2m - 6$: its prefix is $\{0, 2, \dots, m - 1\}$, of defect $m - 3$, and adding $m + 1$ makes exactly the $m - 3$ pairs of difference one fail. The third-gap induction therefore provides

candidate branches, but it does not by itself classify every case of equality in those branches. An efficient next step is to classify endpoint Schur defect two before attempting a larger search.

Question V.166.2 (Bounded Schur defect). *For a fixed integer $s \geq 2$, describe all positive real sets $C = \{c_1 < \dots < c_n\}$ satisfying*

$$\binom{n}{2} - \#\{x + y = z : x, y, z \in C\} = s.$$

For n sufficiently large in terms of s , are all such sets positive real dilates of finitely many integer terminal perturbations of $\{1, \dots, n\}$? If so, obtain effective bounds on the length of the perturbed segment and on the required threshold in n .

[write] **7 October 2026; sources 54–59.** *Answered by source 57* (Section II.291): for $n \geq 5s - 1$, $h(C) = s$ holds exactly for the positive dilates of $\{j + \lambda_j : 1 \leq j \leq n\}$ with (λ_j) a partition of s (nonnegative, nondecreasing, the perturbed positions among the last s), so there are $p(s)$ classes (its Theorem II.292.3); for $s = 2$ the threshold $n = 6$ is optimal (its Theorem II.292.4). The optimal threshold for general s is open (item II.53).

[write] **7 October 2026; sources 60–69.** *Threshold lowered by source 66* for $s \geq 3$: $n \geq 3s + 1$ (its Theorem II.302.1); $n = 2s + 1$ is not enough (its Proposition II.302.4).

The unit-defect proof uses the first imperfect reflection row and the next perfect row. With several defective rows, their locations can interact. A classification that keeps track of those interactions could turn the endpoint recurrence into a systematic description of the top of the energy spectrum.

Question V.166.3 (A larger rigidity window). *Determine sharp edit or containment bounds in terms of $D(A) = (M_m - E(A))/4$ when $D(A) = O(m)$ or $D(A) = o(m^2)$. For example, for each fixed K , does $D(A) \leq Km$ force all but $O_K(1)$ points into a single arithmetic progression whose length is $m + O_K(1)$? Identify any necessary exceptions and the optimal dependence on K .*

[write] **7 October 2026; sources 60–69.** *Advanced by source 66*: above the fourth level, $E(A) \geq M_m - 8m + 20$, every set lies in an $(m + 2)$ -term progression and contains an $(m - 2)$ -term one (its Corollary II.309.1); the linear-defect question is its Question V.249.6.

The theorem here proves exact one-point repair below the third level. Extending this to a fixed multiple of m would connect the finite classification to mesoscopic stability. A statement at this scale would require new arguments; the present induction does not supply uniform control once many endpoint defects are allowed.

Question V.166.4 (Weighted ordered defects). *Is there a useful nonnegative weighted analogue of (F.50.2), with an explicit extremal weighted profile and a rigidity statement at small deficit? Which choices of mass and magnitude constraints preserve an ordered-completion interpretation?*

Uniform indicator weights are essential to the integer-valued defect and its four-unit spacing. A weighted formulation should specify its normalization before guessing a counterpart of M_m ; the support cardinality alone does not control weighted energy.

Question V.166.5 (Effective cyclic rectification). *Find explicit, checkable hypotheses on $A \subset \mathbb{Z}/p\mathbb{Z}$ that allow the first three torsion-free energy levels to be transferred by a Freiman 2-isomorphism. Can a suitable density or sumset condition be combined with a near-extremal energy hypothesis to force those hypotheses, with useful constants?*

The full cyclic group has energy m^3 , so a torsion-free bound cannot hold for arbitrary cyclic sets. The useful objective is an effective rectification theorem that names and excludes its periodic or wraparound obstructions.

V.166.2 Coset envelopes and puncture rigidity

Question V.166.6 (The largest envelope range). *Determine the largest ε_* for which every nonempty finite set in every abelian group with $\varepsilon < \varepsilon_*$ has coset distance at most $\tau(\varepsilon)$. Determine separately the largest range in which the closest coset is unique and is given by the autocorrelation threshold $r(d) > m/2$.*

[write] **7 October 2026; sources 60–69.** *Answered by source 68:* $\varepsilon_* = 1/4$ for existence and the autocorrelation subgroup, optimal even in cyclic groups, and $2/9$ for uniqueness of the closest coset, also optimal (Section II.296).

Our common sufficient range is $\varepsilon < 1/9$. The three properties need not have the same threshold. At $\varepsilon = 1/4$, a two-point set in $\mathbb{Z}/3\mathbb{Z}$ has several closest cosets, so uniqueness cannot be extended through that endpoint without modification. Failure of one step in the present proof is not evidence that the envelope itself fails there.

Question V.166.7 (A sharp pointwise autocorrelation inequality). *With $q(d) = |A \cap (A + d)|/|A|$, does the inequality*

$$q(d)(1 - q(d)) \leq \varepsilon$$

hold universally? If it fails, determine the best constant in $q(d)(1 - q(d)) \leq C\varepsilon$ and classify the limiting examples.

[write] **7 October 2026; sources 60–69.** *Answered by source 68* (Section II.296): yes, $q(d)(1 - q(d)) \leq \varepsilon$ holds in every abelian group, with all equality cases at $0 < r_A(d) < m$ (its Theorems II.297.1 and II.297.7).

The proved inequality has $C = 2$. A punctured coset supplies equality in the proposed factor-one inequality for $d \in H \setminus L$ when $A = H \setminus L$, so a universal constant below one is impossible. A better constant would directly strengthen the subgroup bootstrap, while the remaining occupancy and complement steps would still need to be checked over the wider range.

Question V.166.8 (The maximal range of nested-hole repair). *Determine the largest range of $s = \eta/\delta^3$ and δ in which the finite repair bound*

$$\min_{D \subset C} |A \Delta (C \setminus D)|/m \leq \delta \tau(s)$$

holds, with D ranging over nonempty proper subcosets of the canonical coset C . Can the sufficient condition $s < 1/9$ be expanded to its optimal range? For a fixed ambient group and cardinality, how do subgroup indices and unattainable intermediate values of s strengthen this finite bound?

[write] **7 October 2026; sources 60–69.** *Range answered by source 68:* the bound holds for all $s \leq 1/4$ while $\varepsilon < 1/4$, the full real domain of the inner inverse, and the outer condition cannot be relaxed (Section II.296); the group-sensitive refinement is its Question V.256.2.

Theorem II.185.11 settles the asymptotic leading constant: it is exactly one. More precisely, its multiplier is $\tau(s)/s = 1 + s + O(s^2)$, and nested-hole examples attain the finite bound at $s = (u - 1)/u^2$ for every integer $u \geq 8$. The theorem also classifies all positive equality cases as two nested holes. Its proof applies the basic coset envelope to the missing set, so any suitable enlargement of that basic range propagates to the repair bound. The unresolved issues are the maximal validity range and the best group-sensitive bound between available equality parameters.

Question V.166.9 (Arithmetic restrictions on the envelope). *For a fixed finite group G and cardinality m , determine the best distance bound at every feasible energy. How much can the universal envelope be improved when G has no subgroup index corresponding to an exact puncture extremizer?*

Equality permits only $\varepsilon = (r - 2)/(r - 1)^2$ at positive defect. In a fixed group, r must be an available subgroup index; in $\mathbb{F}_p^n$ it is a power of p . The quantitative hole remainder and integer constraints on $|A \Delta C|$ provide two possible routes to a group-sensitive bound between those equality points.

Question V.166.10 (Stable collision equality and labels). *If coordinate-cube collisions nearly attain the coset-distance envelope, must the domain be close, at a quantitatively sharp rate, to a product with one punctured coordinate? Under an additional near-maximal labelled-collision condition, which labels can occur on such a domain, and how close must they be to compatible additive or multilinear labels?*

The equality argument separates the domain geometry from labels. A stability proof would have to convert a small failure of corner closure into a product approximation and then control the labels on the repaired domain. Domain closeness alone does not imply label compatibility. This is a natural continuation of the collision amplification work in report source 15 [202].

V.166.3 The fourth-uniformity extremal problem

Question V.166.11 (The exact value of c_4^{odd}). *Determine the constant in (F.50.4) for $d = 4$ and all its extremizers or extremizing sequences. Is the supremum attained on a finite group? Can an extremizer always be realized on a cyclic quotient, or is a higher-rank group needed?*

The established interval has width about 0.0880. The five-point function is an exact lower-bound witness, not a proposed solution of the optimization problem. Searches over small groups should retain exact or interval-certified witnesses; a floating-point local optimum does not certify a universal bound or even optimality on that group.

Question V.166.12 (A stronger moment certificate). *Improve the rational certificate in Section III.90 by using more of the joint constraints on $t(h) = (\mathbb{E}_x f(x)f(x+h))^2$ and $w(h) = \|f(\cdot)f(\cdot+h)\|_{U_2}^4 - t(h)^2$. Can a finite, exact sum-of-squares certificate reach the sharp constant, or is information beyond finitely many moments necessary?*

The proof drops several nonnegative Fourier contributions and applies Jensen to a quadratic expression. Optimizing its coefficients alone need not recover the true extremal value. A useful next certificate should record which relaxations can be simultaneously sharp and exploit any incompatibility.

Question V.166.13 (Spectral support and torsion). *Describe the Fourier support of a fourth-uniformity extremizer. What is the minimum possible number of opposite character pairs, and do two pairs suffice? Does order five play an intrinsic extremal role, or does it merely supply a convenient finite witness? Determine how the answer changes when all character orders are restricted to exceed a prescribed bound.*

The single-pair class is rigorously excluded. The dual identity $\mathcal{D}_4\phi_1 = 111\phi_1 + 27\phi_3 + \phi_5$ explains why a third harmonic is forced in an improving perturbation. Subsequent iterations may generate further harmonics or torsion-dependent aliases. Their compatibility is a structural question, not just a coefficient calculation.

Question V.166.14 (Stability and higher dimensions). *Prove a useful stability theorem for near-extremizers of c_4^{odd} , once its extremal models are identified. Determine the constants c_d^{odd} for $d \geq 5$ and whether there are strict gaps between consecutive values.*

Monotonicity supplies $c_d^{\text{odd}} \leq (832/4795)^{1/4}$ for $d \geq 4$, but does not identify any higher-dimensional extremal model. A repeated derivative argument must retain the mean of the derivative, as the proof here does; centering the original function does not center each derivative.

V.166.4 Integration and longer-term consequences

Question V.166.15 (Formal proof of the finite certificates). *Formalize the endpoint energy identity, the unit-Schur-defect lemma, the complement-energy identity, and the rational fourth-uniformity certificate in Lean. Which finite combinatorial identities can be checked by a small trusted kernel, and which analytic transitions need additional finite Fourier infrastructure?*

Appendix V.167 proposes a dependency order and gives the precise normalized-to-unnormalized translation. The shortest formalization route should begin with exact finite sums and rational polynomial identities, then add the structural equality arguments.

Question V.166.16 (Propagation to a global quantitative theorem). *Identify a complete quantitative argument in which one of the local gains here survives all subsequent losses. Does the improved fourth coefficient change a finite-parameter cube or progression threshold of independent interest? Can any such gain be propagated through a full density-increment proof with an explicit final bound?*

A convincing answer should expose every dependency and retain the same conventions for ambient density, uniformity, and degenerate cubes. Improving a supplementary local coefficient is not sufficient evidence of a better global Szemerédi bound.

For near-term work, the most concrete targets are the third-level energy classification, the sharp pointwise autocorrelation constant, and a better exact fourth-uniformity certificate. They each have a specified gap left by the present proofs, a finite test environment, and a plausible route to a further theorem.

V.167 Normalization and a formalization roadmap for ProveIt

Source 45, delivered Appendix A.

V.167.1 Relevant existing declarations

The inspected formal development is `Combinatorics/Ramsey/Lean/GowersSzemerédi/` at the pin in Appendix V.169. The following declarations give the useful interfaces. Their presence is a source-level observation; this report does not claim to have compiled a Lean extension.

File and declaration	Role for this manuscript
<code>Definitions.lean: cubeDifference, gowersNorm, cubeCount</code>	Iterated multiplicative differences, unnormalized Gowers norms, and ordered additive-cube counts
<code>Definitions.lean: UniformOfDegree, UniformSetOfDegree</code>	The parameter α in (III.91.5) is normalized so that degree $d - 1$ uniformity means $\ f\ _{U^d}^{2^d} \leq \alpha$
<code>Proofs03CubeUpper.lean: lemma_3_10_holds</code>	Existing proof of the cube upper bound using Minkowski; a strengthened corollary can be added beside it
<code>Sections06_07.lean: theorem_7_2, proposition_7_3</code>	Statements of the general structural energy reduction; the torsion-free top-level results are separate, stronger conclusions under much stronger energy hypotheses

The paper’s sign convention uses $x + \omega \cdot h$, while the development’s affine cube argument uses a minus sign. Replacing each h_i by $-h_i$ preserves the sum, so this difference has no numerical effect.

V.167.2 The essential scaling identity

Write $\Gamma_d(f)$ for the repository’s unnormalized `gowersNormdf`, and let $N = |G|$. For the finite cyclic instance implemented there and $d \geq 1$,

$$\Gamma_d(f) = N^{(d+1)/2^d} \|f\|_{U^d}, \quad \Gamma_d(\mathbf{1}_A)^{2^d} = \text{cubeCount}_d(A). \quad (\text{V.167.1})$$

The latter equality is proved locally in `Proofs03CubeUpper.lean` as `gowersNorm_indicator_pow`. It counts degenerate cubes as well.

In particular, the real fourth-uniformity theorem becomes the denominator-free inequality

$$4795 \Gamma_2(f)^{16} \leq 832 N^7 \Gamma_4(f)^{16} \quad (\text{V.167.2})$$

for real centered f on an odd cyclic group. Indeed $\|f\|_{U_2}^{16} = N^{-12} \Gamma_2(f)^{16}$ and $\|f\|_{U_4}^{16} = N^{-5} \Gamma_4(f)^{16}$. Omitting the factor N^7 would be a normalization error.

For the finite energy statements, $E(A) = \text{Finset.addEnergy}(A, A)$ under the ordered quadruple convention used by the inspected statements. The torsion-free theorem can be stated first for $\text{Finset } \mathbb{Z}$, with cardinality hypotheses made explicit. Generalizing the ambient group is a later interface step rather than a prerequisite for the first formal proof.

V.167.3 A proposed order of formalization

1. **Ordered finite counting.** Formalize the energy increment on adjoining a greatest element, the bound on its Schur count, and the ordered missing-completion identity. Use integer arithmetic for $4D(A) = M_m - E(A)$, or clear the denominator three in M_m .
2. **Equality and unit defect.** Encode the reflection characterization of each Schur row. Prove zero defect, then the unique defective-row lemma, then the second and third energy results. This order avoids a separate full classification of four-point energy sets.
3. **Autocorrelation and complements.** Establish $\sum_d r(d) = m^2$, $E(A) = \sum_d r(d)^2$, the overlap inequality, and the exact complement expansion. The polynomial remainder can be verified after clearing denominators. Construct the canonical subgroup with a finite-set closure argument before introducing the radical expression τ .
4. **Norm certificates.** Prove the Fourier square identity for U^3 and the opposite-frequency pairing inequality on odd-order groups. The rational quadratic certificate is a separate real polynomial identity; it should not be entangled with Fourier bookkeeping. Then formalize the nonnegative autocorrelation moment expansion and combine the lemmas.
5. **Exact finite examples.** Verify the five-point cube sums by finite enumeration over $\mathbb{Z}/5\mathbb{Z}$ and the cosine sign tables by Boolean enumeration. These have small, fully specified domains and require no numerical approximation.
6. **Cube-counting interface.** Prove the centered quartic expansion, including the odd-order rank argument, or reuse it if subsequently formalized in the repository. Combine with (V.167.2) to obtain (III.91.5) with the existing uniformity parameter.

No declaration names proposed here assert that such theorems have already been added to the repository. The package is a standalone article suitable for later integration under `Research/GowersSzemerédi/rigidity-gaps/`, or for a further source intake into the existing thematic report. If it is consolidated, retain the provenance distinctions and do not silently replace the earlier reports' hypotheses or their historical claims.

V.168 Exact certificates and reproducibility

Source 45, delivered Appendix B.

The package contains three independent standard-library Python scripts. They use integers and `Fractions.Fraction` for all mathematical comparisons. Decimal values in the prose are evaluations of exact expressions and are not inputs to any theorem. The searches cover only the explicitly recorded finite domains. Their role is to expose counting, normalization, equality-case, and transcription errors; they do not replace the universal proofs.

V.168.1 Energy and Schur counts

`code/verify_energy.py` computes energy by counting pair-difference multiplicities and independently counts failed ordered triple completions. It verifies the endpoint recurrence using a reflected Schur count, checks progression equality, the first gap, the complete second-level classification in its stated range, and the third gap. The recorded run covers:

- all 12 911 subsets of $\{0, \dots, 14\}$ containing zero and having cardinality between one and nine;
- all 39 202 subsets of $\{1, \dots, 16\}$ with cardinality between one and eight for the positive Schur counts;
- all 276 single interior holes in $\{0, \dots, m\}$ for $2 \leq m \leq 24$, checking the exact one-hole formula;
- five additional rational or large-gap fixtures, together with explicit small-cardinality and torsion counterexamples to overbroad versions of the statements.

The five largest tested cardinalities give the following observed top three energies, consistent with the proved formula:

m	M_m	$M_m - 4(m - 2)$	$M_m - 8(m - 3)$
5	85	73	69
6	146	130	122
7	231	211	199
8	344	320	304
9	489	461	441

No full third-level equality classification is asserted or tested by this script. The output is `data/energy_checks.json`.

V.168.2 Cosets and coordinate collisions

`code/verify_cosets.py` computes exact energy and enumerates subgroup cosets to check the canonical construction and unique minimum. It verifies the full polynomial remainder after clearing denominators, checks the equality classification, checks puncture repair on its admissible finite cases, and counts coordinate-cube collisions directly in small product groups. The recorded run exhausts 131 053 nonempty cyclic subsets in groups of orders two through sixteen and examines 1 597 product-group subsets. Three further cyclic fixtures are added for the recursive puncture theorem. There are 328 high-energy cases overall, and 320 outer-envelope equality cases, of which 101 have positive defect.

For the nested-hole bound, the script minimizes over all coset holes inside the fixed canonical C in each of 107 admissible concrete cases. Six have positive η : five attain the bound and one is strict. Three lie beyond the smaller preliminary cutoff $\eta/\delta^3 \leq 1/16$. The radical-free comparison is

$$2j \leq d, \quad dj(d - j) \leq J_{\text{int}}, \quad J_{\text{int}} = m^3 - E(A) - dm^2 + d^2m,$$

where $d = |A \triangle C|$ and j is the minimum edit count. An additional 427 nested-hole equality examples and 18 strict outlier examples are checked through exact closed formulas; no large ambient groups are enumerated for these checks. The output is `data/coset_checks.json`. These numbers describe the test suite, not a count of all isomorphism types or a finite proof of the general theorems.

V.168.3 The fourth-uniformity certificate

`code/verify_u4.py` performs three complementary checks. First, it expands the rational quadratic certificate exactly and recovers

$$\frac{4795}{832} + \frac{3}{8} \left(V - \frac{3}{2} \right) + \frac{9}{8} \left(V - \frac{3}{2} \right)^2.$$

Second, it computes the fifteen-vertex dual-cube sign table by direct enumeration of 2^{15} assignments and independently by dynamic programming in the generating polynomial. It also checks the sixteen-vertex modular sign counts for character orders 3, 5, 7, 9, including the order-three phase terms. Third, it enumerates every ordered cube of dimensions two, three, and four on $\mathbb{Z}/5\mathbb{Z}$ for the specified vector $(2, 1, -2, -2, 1)$.

The fourth-dimensional count includes all $5^5 = 3125$ parameter tuples and yields the exact integer product sum 6 468 874. Both the histogram and its weighted sum are retained in `certificates/u4_certificate.json`. The generating identity in the proof, together with the finite recurrence printed there, makes the cosine counts independently reproducible without trusting floating-point software.

V.168.4 Commands and verification boundaries

From the package root, run:

```
python3 code/verify_energy.py --output data/energy_checks.json
python3 code/verify_cosets.py
python3 code/verify_u4.py --json certificates/u4_certificate.json
latexmk -pdf -interaction=nonstopmode -halt-on-error \
  gowers_rigidity_gaps.tex
```

The convenience command `python3code/verify_all.py` runs the three mathematical scripts in sequence. Python 3.10 or later is sufficient. The scripts using assertions explicitly reject optimized Python modes that would disable them. The energy script uses explicit failure checks. No internet access or third-party Python package is needed to reproduce the exact mathematical outputs.

Internal AI cross-checks of the proofs are included in `reviews/`. Different analysis passes checked the energy, coset, fourth-uniformity, and integration arguments. These records make the checks inspectable; they are not external peer-review reports or proof-assistant certificates. The final PDF was compiled, checked for unresolved references and layout warnings, and visually inspected. Build and package details are recorded in `data/build_report.json`.

V.169 Source ledger and limits of the audit

Source 45, delivered Appendix C.

The repository snapshot used for this article is

6ce33539ea468f51719a211c18aabdde262a24dd

from <https://github.com/VladimirReshetnikov/ProveIt>, inspected on 7 October 2026 UTC. All paths below are relative to `Combinatorics/Ramsey/`. Exact blob identifiers are also recorded in the machine-readable companion `data/source_manifest.json`.

1. `Papers/sz-thm-gowers-proof/sz-thm-gowers-proof.tex`. This is the edited repository transcription of the original paper. The primary published paper was checked separately. The principal local targets are Lemma 3.10 and the energy discussion in Section 7 immediately preceding Theorem 7.2 [1].
2. `Research/GowersSzemerédi/local-quantitative-refinements/README.md` and `article.tex` in the same directory. The consolidated article at this pin includes sources 01–14. Later sources 15–23 are represented in the directory by integration notes, manifests, scripts, data, and figures, rather than by a completed further consolidation into that article. Their available descriptions and relevant verification material were inspected for overlap; this was not a line-by-line audit of every earlier proof.

3. `Lean/GowersSzemerédi/Definitions.lean`, `Lean/GowersSzemerédi/Sections06_07.lean`, and `Lean/GowersSzemerédi/Proofs03CubeUpper.lean`. These were read to check definitions and interfaces. No new Lean theorem or compilation result is supplied.

V.169.1 Precise links to the existing research questions

The relevant consolidated labels are given literally so that future integration need not depend on page numbering.

Coset envelope and its range.

Source 14’s theorem `gsr:en:14:thm:energy-sharp-stability` proves the sharp envelope for $\varepsilon \leq 1/100$. Question 4, consolidated label `gsr:fz:q:II:cosetrangle`, asks for a wider range and all exact equality examples. Section II.185 enlarges the range to $\varepsilon < 1/9$ and classifies every equality case in that range. It does not settle the largest possible range.

Coordinate collisions.

Source 14’s `gsr:en:14:thm:collision-sharp-stability` has the corresponding small-range product-coset conclusion. Its sharpness examples and exact endpoint results do not classify all positive-defect envelope equality domains. Source 15’s collision amplification material still uses $1/100$ in its verifier. The new range and punctured-coordinate equality classification are proved in Section II.185.

Higher real norm gaps.

The higher-dimensional question occurs in source 03 Question 1, source 04 Question 3, and source 14 Question 11, consolidated label `gsr:fz:q:III:normgap`. Section III.90 improves the $d = 4$ bounds and rules out the single-character model. The exact constant remains unresolved.

The fourth-order cube term.

Source 04’s proposition `gsr:cc:04:cub:quartic`, with corresponding results in sources 03, 08, and 09, gives the coefficient P_d . The real U^3/U^2 gap is also an existing result. Section III.91 credits these facts and substitutes the new fourth-uniformity constant.

V.169.2 Published context and priority

The primary-source comparison included Gowers’s original paper, the inverse Pollard theorem of Nazarewicz–O’Brien–O’Neill–Staples, Hegyvári’s paper on the additive-energy spectrum, Shao’s large-energy structural theorem, and Eisner–Tao’s near-extremal Gowers-norm work [1, 190, 189, 191, 7]. The first positive torsion-free energy gap is already implied by the inverse Pollard result. The endpoint recurrence, congruence, and displaced-endpoint construction are also credited to prior work. An independent elementary proof is included because of its role in the new finite classification.

No primary source inspected was found to contain the exact second-level classification, the third exclusion interval, the expanded coset range with its complete puncture equality theorem, or the fourth-uniformity constant proved here. Absence from this search is not a proof of global novelty. Before a publication claims priority, an expert review should check the relevant inverse sumset, rearrangement, and extremal-uniformity literature more broadly. The mathematical claims of the manuscript are its explicit statements with their proofs; the priority claim is deliberately narrower.

V.170 Algorithms, reproducibility, and proof status

Source 46, delivered Section 10.

V.170.1 An effective construction pipeline

Given the coefficients of a symmetric binary trilinear T in a basis, the proof is algorithmic. Compute the matrix $C_{ij} = T(e_i, e_i, e_j) + T(e_i, e_j, e_j)$. A symplectic elimination yields its rank, a maximal isotropic subspace, and a decomposition into elementary wedges. These determine a canonical correction S after choices of pivot ordering have been fixed. Then (IV.29.10) integrates $R = T - S$, and $e(P)$ is an extremizer.

The matrix calculations use ordinary finite-field Gaussian elimination, with cubic arithmetic complexity in the dimension. Once the trilinear coefficient table has been formed, subtracting the elementary corrections and writing the polynomial also takes polynomially many operations. The output can be kept symbolically with $O(n^3)$ coefficients; it need not be expanded into its 2^n -entry value table. A naive coefficientwise implementation of all corrections uses $O(n^4)$ field operations; this is an upper bound, not an optimal-complexity claim. For trace tensors the closed rank formula avoids constructing the matrix if only the extremal value or maximal restriction dimension is required.

For general p at degree $p + 1$, the same obstruction matrix and symplectic elimination construct an optimal elementary repair and an optimal-dimensional integrable restriction. Writing the full symmetric tensor in coordinates has a size depending on p and n ; no dimension-free runtime is claimed for an input that is already a high-order coefficient table.

V.170.2 Exact finite checks supplied with this paper

The accompanying `verify.py` uses only the Python standard library. It performs finite-field arithmetic in polynomial quotient representations, integer calculations, and rational comparisons. It does not use floating point to validate any theorem. The deterministic seed is 172917. The supplied `validation.json` records the field moduli, coefficients, phase numerators, witnesses, counts, and results.

Check family	Number checked
Scalar leading differences and phase orders, $p = 2, 3, 5, 7, 1 \leq d \leq 20$	80
Explicit trace integrators: all listed direction tuples and basepoints	141 781
Nonzero trace coefficients: radical, isotropic, and wedge certificates	770
Explicit repeated-variable nonintegrability witnesses	5
Fourth-root single-function tables for binary spectral moments	692
General symmetric binary cubic extremizers: cube/basepoint checks	209 664
Mixed eight-function tests for general symmetric binary tensors	36
Seven-function tests for general symmetric binary tensors	36

The trace-coefficient checks cover every nonzero coefficient in $\mathbb{F}_{2^n}$ for $1 \leq n \leq 8$, in $\mathbb{F}_{3^n}$ for $1 \leq n \leq 4$, and in $\mathbb{F}_{5^n}$ for $1 \leq n \leq 3$. The general binary cubic checks use nine coefficient tables in dimensions two, three, and four. They explicitly construct an extremizing phase and check its total derivative against the residual tensor at every basepoint and every direction triple. The script also verifies the predicted Walsh spectral magnitudes and the exact extremal moments.

For the two specified coefficients over $\mathbb{F}_4$, all $4^4 = 256$ fourth-root function tables are exhausted. For $c = 1$, their maximum moment is 1; for $c = \theta$, it is $7/16$, whereas the unrestricted maximum proved in this paper is $1/2$. The latter value is attained using eighth roots. The larger-field function-table checks are sampled rather than exhaustive; the JSON explicitly distinguishes these cases. Their largest observed values must not be mistaken for the proved optima.

Run the checks with

```
python verify.py --output validation.json
```

The program raises an exception on a failed assertion and writes a final **PASS** status only after completing the checks. These computations are useful protection against normalization, sign, and phase-depth errors.

They are *not* the proof of any universal theorem, and they are not a Lean kernel check. The general proofs are those in the preceding sections.

V.170.3 Repository integration and formalization boundary

A suggested destination is a new, descriptively named subdirectory under `Combinatorics/Ramsey/Research/GowersSzemerédi/`; this package does not modify any existing repository file. The source audit records the inspected blob identifiers of the Gowers transcription and the refinement inventory. Blob identifiers should not be mistaken for commit identifiers.

The most economical formalization order is to start with the finite binary core: alternating obstruction; quadratic Walsh spectrum; the cube change of variables; exact upper bound; explicit residual integrator; sharpness; and the seven-function corollary. The trace rank calculation and the general- p coordinate construction form subsequent independent modules. The accompanying `FORMALIZATION.md` spells out proposed statements and dependencies. It is a roadmap, not purported compiled Lean code.

V.171 Further research questions and concrete next targets

Source 46, delivered Section 11.

The questions below are proposed continuations of this work. Calling them questions means they are not resolved here; it is not a claim that each has never been considered elsewhere.

Research question V.171.1 (The exact seven-function norm). For a symmetric binary trilinear T , determine the exact supremum of $|\Gamma_T|$ over seven 1-bounded functions. Does it depend only on rank C_T , as the eight-function norm does? Is the exponent $r/4$ in Theorem IV.28.1 optimal?

The cube estimate is sharp, but equality must also survive the initial Cauchy–Schwarz step. An exact analysis of that compatibility would decide whether the codimension constant 2 can be decreased. Small-rank blocks and tensor products provide explicit test cases.

Research question V.171.2 (Classification of all cubic extremizers). Describe all f for which $M_T(f) = 2^{-r/2}$, modulo natural phase multiplications and linear changes of coordinates.

Proposition IV.28.4 says that every transformed four-point function must lie in the positive spectral subspace, and that $|f| = 1$. The missing step is to solve these simultaneous compatibility equations for f . The constructed nonclassical phases give one family, not an exhaustive classification.

Research question V.171.3 (From spectral to structural stability). If $M_T(f) \geq (1 - \varepsilon)2^{-r/2}$, can one approximate f in normalized L^2 by an exact extremizer, with a dimension-independent error tending to zero with ε ?

The exact defect identity is a useful starting point because it has no unspecified constants. The obstacle is that the favorable eigenspaces vary with (A, B) ; separate projection estimates do not immediately produce one globally consistent phase.

Research question V.171.4 (The ordinary slice-rank repair gap). Determine $\rho(T)$ within the interval $r/2 \leq \rho(T) \leq r$. Is $\rho(T) = r$ for all symmetric binary trilinear forms, or can several obstruction blocks be repaired more economically together?

The rank-two case is settled in Proposition IV.32.4. In larger rank, a sum of nonsymmetric slices can be symmetric without each summand being symmetric. This is exactly why the elementary repair theorem does not answer the ordinary slice-rank question.

Research question V.171.5 (Higher phase-alphabet constraints). Determine the maximum of $M_T(f)$ when f is restricted to a prescribed finite root-of-unity alphabet, especially fourth roots rather than eighth roots in characteristic two.

The exhaustive $\mathbb{F}_4$ calculation exhibits a strict gap $7/16 < 1/2$. A general theorem should describe which residual diagonal coefficients force this gap and whether product constructions preserve an exact alphabet-dependent norm formula.

Research question V.171.6 (Odd-characteristic analogues). At degree $p + 1$, is there a sharp mixed cube inequality controlled by the rank of C_T over $\mathbb{F}_p$, and can its extremizers be constructed by symplectic blocks?

The algebraic restriction and repair statements already hold for every prime. The analytic proof in Section IV.27 exploits a special binary reparametrization of the eight vertices and a quadratic kernel. An odd-characteristic replacement would require a different analytic mechanism, not just changing signs in the present proof.

Research question V.171.7 (Higher-degree trace restrictions). For $d \geq p + 2$, determine the largest subspace $W \leq \mathbb{F}_{p^n}$ on which $T_{c,d}$ integrates.

Theorem IV.31.1 excludes integration on the whole extension field. On a subspace, the repeated-variable constraints involve all products of the extra directions from W . This becomes a simultaneous isotropy problem for a family of forms C_{cz} , rather than one alternating form. Product-set structure of W should therefore become relevant.

Research question V.171.8 (Robustness before exact symmetry). Can a quantitatively approximate symmetric tensor be brought into the framework of Theorem IV.27.1 with a loss that improves a complete low-characteristic symmetrization pipeline?

A useful result must separately track the error of symmetrization, the rank of the repeated-variable obstruction, and the cost of restriction or repair. The sharp symmetric estimate alone does not justify silently replacing constants at all those earlier stages.

Research question V.171.9 (Frobenius-twisted trace families). Extend the exact obstruction-rank calculation to symmetric forms built from sums of trace monomials involving Frobenius powers in the arguments.

The radical will again be described by a linearized polynomial, but the number of roots need not have the two-regime form of (IV.32.1). Explicit families could provide a broader supply of sharp cubic examples and nontrivial restriction geometries.

Research question V.171.10 (Formal certificates and algorithmic output). Construct machine-checkable certificates for the symplectic decomposition, residual integrator, and extremal value, and formalize their verification in Lean with finite normalized sums.

A certificate-based implementation can separate potentially elaborate search from a short verifier. The package's exact arithmetic checks are prototypes of the data to certify; they are not a substitute for the kernel-verified implications.

V.172 Conclusion

Source 46, delivered Section 12.

The factorial restriction in polarization is not one uniform obstruction. On the scalar prime field it can always be repaired, at exactly the phase order $p^{\lceil d/(p-1) \rceil}$. For extension-field trace tensors it can be a genuine

obstruction to every circle-valued integrator. At the first obstructed degree, an alternating form measures the exact geometric cost of restriction and the exact number of elementary symmetric repairs.

For binary cubic tensors, that same form has an exact analytic meaning: the twisted mixed cube norm is $2 - \text{rank } C_T/2$. The quadratic convolution proof and the matching nonclassical-phase construction yield sharp agreement thresholds and an explicit stability identity. They also improve the symmetric seven-function restriction bound from $8 \log_2(1/\delta)$ to $2 \log_2(1/\delta)$. These are local, fully specified advances at a genuine structural bottleneck of the Gowers method, rather than an unverified claim of a new global Szemerédi bound.

V.173 Proof dependency map and scope checks

Source 46, delivered Appendix A.

The main analytic theorem has the following dependency chain:

binary cube change of variables $\longrightarrow$ quadratic Walsh spectrum $\longrightarrow$ mixed upper bound,
alternating-form decomposition + explicit binary cubic integration $\longrightarrow$ matching extremizer.

Neither chain uses an inverse theorem. The seven-function corollary adds one application of Cauchy–Schwarz and the maximal-isotropic dimension formula. The trace applications add only Frobenius adjointness, nondegeneracy of the trace pairing, and root counts in the cyclic group $\mathbb{F}_q^\times$.

The most consequential scope checks are these. Symmetry of T is essential in the sharp cubic proof. Tensor degree three and characteristic two are essential to its quadratic-kernel reduction. “Integrates” permits nonclassical circle-valued polynomials unless “classical” is explicitly specified. The exact repair number refers to the elementary family (IV.32.5), whereas ordinary slice-rank repair only has the bounds in Proposition IV.32.4. The exact seven-function norm, the complete extremizer classification, and an improved global progression-free-set bound are not proved in this article.

V.174 Relation to the Gowers argument and the formal project

Source 47, delivered Section 10.

V.174.1 A precise local generalization

The relevant source stage is Gowers’s Lemma 17.1 and Proposition 17.2, at printed pages 577–578 [1]. The source uses finite differences and a factorial denominator to turn structured frequencies into polynomial phase information. Our result treats a Boolean vector-space analogue in which that denominator is unavailable, and quantifies the additional obstruction for an already supplied symmetric tensor.

The inspected formal statement in `Lean/GowersSzemerédi/Sections17_18.lean` explicitly assumes a prime cyclic modulus and invertibility of $(k+1)!$, and concludes a phase of degree at most $k+1$. Its source comment records the correction of the printed degree. In contrast, our d is the degree of a genuinely multilinear tensor including its final evaluation slot, so $d = k+1$ at the selected-energy interface. Dividing the project’s unnormalized Fourier-square sum by $|V|^{d+1}$ gives our normalization.

There are three substantive differences, not just notation. The ambient space is $\mathbb{F}_2^n$ rather than a large prime cyclic group. The phase takes values in $\mathbb{R}/\mathbb{Z}$ rather than being required to be a classical polynomial with values in the base field. Finally, the sharp criterion assumes full symmetry and high selected energy for a specified tensor; it is not a procedure that finds that tensor from an arbitrary positive uniformity norm.

V.174.2 Multiaffine terms must not be silently discarded

The project’s `Definitions.lean` defines `IsMultilinear` as a sum of square-free monomials, including constant and lower-degree terms. It is therefore a multiaffine predicate in the usual terminology. Our tensor is homogeneous multilinear in every slot. A general multiaffine frequency selector cannot be covered by simply erasing its lower-order terms. Such a reduction would need its own phase correction or energy estimate, and no such reduction is assumed here.

Similarly, all translation averages in the proof range over the full vector space. On a progression, a Bohr neighborhood, or a nonuniformly weighted set, a translation may change the averaging measure. The exact period identity still holds pointwise, but its Fourier-cancellation and projective-averaging consequences acquire boundary or measure errors. A new global quantitative Szemerédi estimate would require those steps, selector construction, and a full downstream loss calculation.

V.174.3 Suggested formalization order

The new elementary results admit a fairly short dependency chain once a finite Boolean Fourier interface is available:

$$\begin{aligned} \text{cube product} &\longrightarrow \text{even periods} \longrightarrow \text{forbidden Fourier coefficients} \\ &\longrightarrow \text{fixed-derivative cap} \longrightarrow \text{canonical induction.} \end{aligned}$$

For the universal theorem, a second chain is

$$\begin{aligned} &\text{pure-map lemmas} \longrightarrow \text{trilinear census} \\ &\longrightarrow \text{support lifting} \longrightarrow \text{defect dichotomy.} \end{aligned}$$

The bridge between the two is the shear rank profile. It is the heaviest background component because it uses finite-dimensional complex spectral theory or an equivalent explicit matrix argument.

The companion `FORMALIZATION.md` records proposed theorem names, inputs, and dependencies. An appropriate implementation should represent characters as actual linear maps $V \rightarrow \mathbb{F}_2$, use normalized finite sums, and prove the cube-product period identity before introducing any division. The sixteen-row census can be a transparent exhaustive Boolean calculation. No unchecked axiom, `sorry`, or proposition-valued declaration in this package is advertised as a formal proof; no Lean files are supplied.

[write] **7 October 2026; sources 47–53.** *Checked at HEAD 06b30ca52.* `proposition_17_2` (`Sections17_18.lean`, line 48) assumes a prime modulus, $k + 1 < N$ and the invertibility of $(k + 1)!$, and concludes a phase of degree at most $k + 1$, as described above; `IsMultilinear` is the multiaffine predicate (`Definitions.lean`, blob unchanged since the pin). The exact companions are `proposition_17_2_holds` (`Proofs17PhaseRemoval.lean`, line 373) and `lemma_17_1_holds` (`Proofs17Phase.lean`, line 374). No statement of this source is formalized; its results are over $\mathbb{F}_2^n$, outside the formal setting.

V.175 Exact verification and source provenance

Source 47, delivered Section 11.

V.175.1 What the finite checks establish

The companion `code/verify.py` uses only the Python standard library and performs no network requests. All its mathematical checks are exact; there is no floating-point optimizer or tolerance-dependent claim in the shipped verifier. It runs safely with or without Python’s assertion optimization, because failures use explicit checks.

The generated record `data/verification.json` includes:

1. 31 644 comparisons of formal cube monomials for the period and conjugation identities. It records separate exponents for $f(x)$ and $\overline{f(x)}$, so these checks also cover zeros and repeated vertices.
2. Exact forbidden-frequency pair cancellations; exhaustive canonical constant counts in degrees 3 through 10; and the rational induction recurrence through degree 1000. An additional 7956 top-derivative checks and 4200 degree checks verify explicit primitives by integer differences modulo 2^d , in dimensions one through four and degrees two through eight, with exhaustive small cases and seeded larger cases.
3. All 16 scalar and all 256 $\mathbb{F}_2^2$ -valued binary trilinear maps, plus all 1024 scalar symmetric trilinear forms on $\mathbb{F}_2^3$. The last family has minimum non-pure support $112/512 = 7/32$.
4. Every nonintegrable symmetric binary tensor in degrees 3 through 6, tested at the constant function. The respective tensor counts are 8, 24, 56, 120, and the maximum constant energies are $1/2, 11/16, 13/16, 57/64$.
5. Exact Gaussian-rational function tests, including zeros and non-unit amplitudes, for canonical energies, fixed-derivative caps, quarter-slice bounds, and independent selected-energy versus full-cube identities in the smallest cases.

The finite tests do not optimize over all bounded functions and do not prove the all-dimensional theorems. Those assertions follow from the mathematical proofs. The census in Lemma IV.48.6 is small enough that the complete proof is printed in the article rather than hidden in a machine certificate.

V.175.2 Pinned interfaces and comparison materials

The observed repository branch ref during the final source inspection was

5c9a442d263632fdcd4e9690154c12c2dcc70b53.

The two interface files read have Git blobs

```
Sections17_18.lean:
18495f50475835ca1c497d85b0052215f865b8ca

Definitions.lean:
97113b7afa6925a2dd4b76641eeaeff09597ab6a.
```

These identify the actual inspected content; no repository-wide build was inferred from statement-file comments. The prior comparison manuscript was read from the user's Library as `boolean_phase_integration.tex`, whose own comparison pin was 797aa3cfe967. Its sections on the cubic law, the shear profile, Boolean integration, and the quartic question were inspected directly.

The original Gowers page image for Proposition 17.2 and Tidor's primary paper were also consulted. The source audit records what was read and distinguishes existing results from the present refinements. No claim of an exhaustive external literature search or an independent referee report is made.

V.176 Further research questions and topics

Source 47, delivered Section 12.

The following are proposed continuations of the results proved here. They are not silently used as hypotheses elsewhere, and they are not presented as a catalogue of already published open problems.

Research question V.176.1 (The septic threshold). Is $M_{7,n}^{\text{sym}} = 15/16$ for every $n \geq 2$? The present proof leaves only

$$\frac{120}{128} \leq M_{7,n}^{\text{sym}} \leq \frac{121}{128}.$$

An improvement must exploit either a better non-pure four-linear defect estimate or simultaneous derivative constraints beyond the scalar rank profile. A support improvement alone should be checked against the compatibility identities before it is used as an energy theorem.

[write] **7 October 2026; sources 60–69.** *Answered by sources 65 and 69:* 15/16 (Section IV.51).

Research question V.176.2 (The all-degree universal extremum). Does the canonical formula $1 - (d + 1)/2^d$ give $M_{d,n}^{\text{sym}}$ in every degree, or is there a first degree and a noncanonical tensor with larger energy? The answer is established here only through degree six. The canonical formula itself is proved in every degree and therefore gives a fixed, dimension-free benchmark for either outcome.

[write] **7 October 2026; sources 60–69.** *Answered by sources 65 and 69:* yes, in every degree $d \geq 3$ and dimension $n \geq 2$ (Section IV.51).

Research question V.176.3 (Optimal non-pure support). Determine the least support of a non-pure symmetric s -linear map for $s \geq 4$, first for scalar-valued maps and then for vector-valued maps. The projection lemma reduces lower bounds in the latter problem to the former. The example $u^{\otimes s} + v^{\otimes s}$ has support $2^{1-s} - 2^{1-2s}$, whereas our uniform lower bound is $7/2^{s+2}$. Are there other support-minimizing families, and what are the stability statements near the minimum?

Research question V.176.4 (Endpoint tensors and phases). Classify the tensors and functions attaining the quintic and sextic thresholds. In the noncanonical branch, equality requires both minimal defect support and equality in every relevant projective spectral bound. The sharpness of the support constant does not by itself prove that all these equalities can occur for derivatives of a single function. For the canonical branch, turn the deficit recursion into a global phase-stability theorem with dimension-independent constants.

[write] **7 October 2026; sources 60–69.** *The tensor part is answered by sources 65 and 69:* the endpoint tensors are exactly the canonical ones modulo integrable tensors, in every degree (Section IV.51); the endpoint functions remain open (65 Research question V.244.3, item IV.11).

Research question V.176.5 (Stability under adjoining inactive coordinates). For which symmetric tensors T_0 does pullback along $V \rightarrow V_0$ preserve $\mathcal{M}(T_0)$? Theorem F.52.1 settles the canonical family. Theorem IV.47.5 reduces every one-dimensional defect image to a binary tensor modulo an integrable correction, but does not settle the remaining optimization. A finite classification of those binary obstruction types, followed by derivative-compatible bounds, is a concrete next target.

Research question V.176.6 (Direct sums and mixed functions). Does the maximum selected energy multiply under direct sums of canonical tensors in degrees at least four? Product functions give a lower bound, not an upper bound. Similarly, determine what replaces the sharp thresholds when different functions occupy the cube vertices. The common function is essential to the exact even-period identity; a mixed theorem requires additional compatibility or a different inequality.

Research question V.176.7 (Odd characteristic and repeated-direction cycles). Find the analogue of the forbidden-frequency argument in characteristic $p > 2$, taking account of the p -repetition transfer conditions in nonclassical integration. The Boolean argument uses that one direction conjugates a derivative and two directions restore it. A useful generalization would retain the full translation action on a repeated-direction product, rather than impose a Boolean period identity that is no longer true.

Research question V.176.8 (Local domains and a complete inverse step). Quantify the period and projective-averaging errors on Bohr neighborhoods or arithmetic progressions, and propagate them through a genuine selector construction and phase-removal argument. An end-to-end quantitative gain must include the cost of producing a symmetric homogeneous tensor from local or multiaffine data; the high-energy certificate alone does not perform this construction.

Research question V.176.9 (Formalization and certified finite models). Formalize the even-period and fixed-derivative lemmas first, then the canonical induction. A separate finite-linear-algebra module can certify the pure-defect classification and the support census. Can the spectral profile be implemented with an explicit finite matrix decomposition that is convenient for both Lean and independently checkable certificates, without obscuring the normalization or using unproved analytic interfaces?

V.177 Conclusion

Source 47, delivered Section 13.

The decisive constraint is not that a tensor happens to use two coordinates. It is that a sliced test function is a derivative, and its further derivatives have exact even-direction periods. Preserving that constraint gives the all-degree canonical extremum and closes the prior quartic stabilization problem.

The universal conclusion requires a second ingredient: a noncanonical symmetric defect cannot be arbitrarily sparse. Pure defects are exactly the canonical obstruction class; non-pure trilinear maps have sharp support $7/32$, and contraction propagates that information to higher degrees. This yields exact symmetric integration thresholds through degree six and the improved uniform bound beyond it. The remaining gaps are explicit: the septic interval has width $1/128$, endpoint phase stability is not classified, and no local-to-global density-increment estimate has been propagated from these results.

V.178 Reproduction and package layout

Source 47, delivered Appendix A.

The archive contains the standalone source `article.tex`, its compiled PDF, a `Makefile`, the verifier and generated data, and source and formalization notes. From the package directory run

```
python3 code/verify.py --max-degree 10
latexmk -pdf -interaction=nonstopmode -halt-on-error article.tex
```

The code uses Python 3.10 or later and the standard library. It also runs under `python3-0`; checks do not rely on assertions. The TeX source uses ordinary pdfLaTeX packages, an embedded bibliography, no external figures, and no shell escape. The delivered PDF is the compiled version of this exact source. No third-party papers, font files, or prior manuscripts are redistributed.

V.179 Worked instances and reproducible exact checks

Source 48, delivered Section 10.

V.179.1 A nontrivial exact-data filter certificate

The bundled example `examples/fejer.json` uses

$$p = 5, \quad B = \{0, 1, 2, 3, 4\}, \quad (\phi(0), \dots, \phi(4)) = (0, 0, 1, 1, 0), \quad s = 2, \quad \eta = 1/3.$$

The full domain has $\mathcal{T}_2(B) = 125$ and $\mathcal{R}_2(B, \phi) = 65$, so its respected proportion $13/25$ is below the target $2/3$. We deliberately use a small exact-data kernel, $L = 3$ and $Q = 2048$, not the worst-case theorem's prescribed value of L . The certified rational table is

$$(q_0, q_1, q_2, q_3, q_4) = \left(1, \frac{595}{2048}, \frac{43}{1024}, \frac{43}{1024}, \frac{595}{2048}\right).$$

Searching all 125 triples gives the exact family mean

$$\frac{2868235983031}{2748779069440} > 0.$$

The best triple, under the specified enumeration order, is $(u, v, c) = (0, 2, 0)$. Its probabilities are $(1, 1, 43/1024, 43/1024, 1)$, with conditional surplus $4377059/786432$. Rounding removes labels 2 and 3 and produces

$$B' = \{0, 1, 4\}, \quad \mathcal{T}_2(B') = \mathcal{R}_2(B', \phi) = 19, \quad F_{1/3}(B') = 19/3.$$

Every number in this paragraph is reproduced using rational arithmetic in `data/fejer_example.json`. The example illustrates why exact certification is useful below the analytic threshold; it is not evidence for changing that threshold without a new proof.

V.179.2 A prescribed-size restriction that is not exact homomorphism

The example `examples/fixed_size.json` uses $p = 7$, $B = \{0, 1, 2, 3, 4\}$, image values $(0, 1, 4, 2, 2)$, order $s = 3$, error $\eta = 1/4$, and prescribed size three. The uniform three-subset average of $F_{1/4}$ is $-15/4$. The deterministic run nevertheless finds

$$B' = \{1, 3, 4\}, \quad \mathcal{T}_3(B') = 111, \quad \mathcal{R}_3(B', \phi) = 93.$$

Thus $F_{1/4}(B') = 39/4 > 0$ and the respected proportion is $31/37 > 3/4$. The algorithm is allowed to exceed a negative initial certificate; the theorem only guarantees nondecrease. The set has the required size exactly, and its respected proportion is not one.

V.179.3 A positive Bernoulli certificate with an affine core

The file `examples/affine_core.json` supplies an example with $p = 17$, $s = 3$, $\eta = 1/4$, and

$$B = \{0, \dots, 7\}, \quad (\phi(0), \dots, \phi(7)) = (0, 1, 2, 3, 4, 5, 12, 15).$$

The first six inclusion probabilities equal $3/4$; the other two equal $1/20$. The exact initial surplus is positive:

$$\mathcal{J}(\mathbf{p}) = \frac{3684947}{12800} > 0.$$

The algorithm retains $\{0, \dots, 5\}$ and returns

$$\mathcal{T}_3(B') = \mathcal{R}_3(B', \phi) = 4332, \quad F_{1/4}(B') = 1083.$$

This is a correctness example for a supplied filter, not a claim that finding such a filter is free. The Fejér example in Section V.179.1 includes the full finite-family search.

V.179.4 What was tested

The program `code/verify.py` uses independent tuple enumeration for small partial coefficients, independent subset enumeration for Bernoulli expectations, and deterministic convolution for larger-order final energies. It does not verify an update by comparing it only with the same update implemented a second time.

Check	Number
Bernoulli inventory instances in cyclic and product groups	48
Individual exact group coefficients compared with tuple enumeration	1160
Exact derivative and replacement comparisons in those instances	144
Signed graph-rounding instances over moduli 2, 3, 4	55
Rounding decisions in those graph instances	185
Distinct-support profiles compared with tuple enumeration	25
Prescribed-size runs, testing every cardinality of each instance	110
Conditional steps in those prescribed-size runs	390
General-system instances, including composite moduli	40
General-system derivative and replacement comparisons	97
Certified rational kernel tables / entries	12 / 78
Exact probability-perturbation comparisons	30
Additional Bernoulli cases at orders 4, 6, 8	9
Additional prescribed-size cases at order 8	4
Signed general-system rounding cases	16
Empty-input and trivial-group regression cases	8

The seed is 20261006. Full results, including the tested moduli and orders, are in `data/verification_results.json`. These checks passed in the delivered environment. Their elapsed time is recorded for reproducibility, not offered as a hardware-independent performance claim.

The general-system tests allow zero and nonunit coefficients, repeated group images, and rational probabilities at endpoints. A separate order-eight check is included because the source application concerns sixteen-tuples. The tests are finite checks of the reference code and identities. They do not replace the mathematical proofs, and no Lean kernel verification is represented by these results.

V.179.5 How to run the code

The implementation needs only the Python standard library and exact `Fraction` arithmetic. From the package root:

```
python code/verify.py --output data/verification_results.json
python code/energy_inventory.py examples/bernoulli.json
python code/energy_inventory.py examples/fixed_size.json
python code/fejer_search.py examples/fejer.json
```

Probabilities and objective coefficients should be integers or rational strings, for example `"3/7"`; inexact floating-point probabilities are rejected. The command-line examples use scalar cyclic groups. The Python API exposes products of cyclic groups and general-system inventories. The complete translated filter search is intentionally streaming and exact, but still has the p^3 parameter factor proved above.

V.180 Integration into the Ramsey development

Source 48, delivered Section 11.

V.180.1 Inspected source interfaces

The relevant repository snapshot is

```
0b4890a1cd1b5417bcf0724db27e35b77bfa96bc.
```

The inspected `Proofs09Restriction.lean` imports `Proofs09Moments.lean` and `Proofs15Restriction.lean`, identifies height-zero arrangements with ordinary additive tuples, and describes its role as proving Lemma 9.3 and Corollary 9.4 [209]. The present results can be developed as an algorithmic layer beside those existence results. They do not require changing the meaning of respected relations or excluding repeated entries.

The companion research manuscript was retrieved as `gowers_polynomial_restriction.tex`. Its source audit refers to snapshot `dca9f06383ccdfc2ed208ccb7c9882214fdb09ac`. That companion is cited for its mathematical results and explicit algorithmic question; its availability as a prior research file should not be confused with a claim that every one of its statements has already been merged or formalized in the currently inspected repository.

A suitable new research location is

```
Combinatorics/Ramsey/Research/GowersSzemerédi/
collision-aware-derandomization/.
```

No repository files were modified for this delivery. The theorem labels in this standalone source use the prefix `gcr:` to reduce collisions if the article is later included in a compilation.

[write] 7 October 2026; sources 47–53. The research location proposed above was not used: by the placement (f8bc5e2ec) this source is source 48 of this report, its labels carry the prefix `gsr:en:48:` or `gsr:fz:48:` in front of `gcr:`, and its files are shipped under the prefix `48-derandomization-`. `Proofs09Restriction.lean` has the same blob at HEAD 06b30ca52 as at 0b4890a1c.

V.180.2 A formalization route without formal exponentials

The implementation and proofs need not formalize analytic exponentials inside a group algebra. A smaller route defines arrays first and proves their tuple interpretation by finite combinatorics.

First define $N_{a,b}$ on labelled finite sets, including the empty tuple and repeated labels. Prove that its Bernoulli expectation is a sum over distinct supports. Next define array insertion by (II.188.4) and prove the inventory invariant by splitting a tuple according to the positions occupied by the newly inserted label. The binomial coefficients then arise from choosing those positions.

For replacement, formalize the Stirling recurrence and the truncated identity

$$(1 + p(e^t - 1))H_p(t) = e^t - 1 \quad \text{through degree } 2s$$

as an identity of coefficient sequences, not analytic functions. This supplies (II.189.3). The rounding theorem then needs only a finite sequence of rational inequalities.

For prescribed sizes, use an integer support-profile array and prove (II.191.2) by counting subsets containing a fixed support. The conditional two-branch averaging identity is a finite binomial identity. Finally, prove that the executable array implementation refines the abstract update. This separation keeps finite algebra, probability, execution, and analytic filter estimates in distinct modules.

The suggested module names in `FORMALIZATION.md` are prospective; no corresponding Lean files are claimed to exist or compile. In particular, this package contains no unverified proof scripts presented as finished formal proofs.

V.180.3 A precise claim ledger

Claim or ingredient	Status
Gowers’s restriction mechanism	Prior source; original PDF pages 516–517 were visually inspected.
Fejér polynomial retention	Prior companion result; needed argument reproved here, with its arithmetic slack retained.
Conditional-expectation principle	Classical; not asserted as new.
Exact collision-aware inventory and response	Proved here for arbitrary finite abelian groups and implemented.
Ambient-size rounding bound	Proved from explicit translation updates; arithmetic and bit costs distinguished.
Prescribed-size theorem	Proved through support profiles and hypergeometric evaluation; all quotas, including endpoints, covered.
General systems and type compression	Proved for the specified finite linear-system model; dependence on ambient state dimension is explicit.
Deterministic source theorem	Complete rational construction preserving the companion’s retention bound; full phase-search cost counted.
Computation	Exact finite tests passed; supplementary trigonometric floating checks are labelled as such.
Formal verification	Not performed for the new results.

Claim or ingredient	Status
Global bounds and priority	No new global Szemerédi bound or exhaustive literature-priority claim.

V.181 Further research questions

Source 48, delivered Section 12.

The results above settle an ambient-size algorithmic interface. The following questions identify improvements that would change its scope, not merely repeat the existence statement.

Research question V.181.1 (Eliminating exhaustive phase search). Can one find a translated Fejér parameter triple with positive surplus in $\tilde{O}_s(np^2)$ or $\tilde{O}_s(p^3)$ time, while retaining Theorem II.194.1’s guarantee? The present total cost is $O_s(np^5)$ because each of p^3 triples is evaluated separately. A useful advance must compute a joint phase objective or a certified conditional estimator, not assume that a random successful triple has large probability merely because the mean surplus is positive.

Research question V.181.2 (Subquadratic ambient storage for graph data). For a scalar map on $\mathbb{F}_p$, the graph inventory uses p^2 locations per coefficient. Can a representation depending on the actual graph sumsets, Fourier support, or low-rank structure preserve exact response probes with substantially less memory? Sparse storage alone is not enough: lower-degree sums can become dense, and subsequent deletions introduce cancellation. Both a running-time bound and a support-growth theorem are needed.

Research question V.181.3 (Sparse and high-dimensional complexity). When ϕ takes values in a high-dimensional vector space, ambient storage can be exponential in its input dimension. Find hypotheses under which the construction is polynomial in the explicit list size and coordinate bit length, rather than in $|G||H|$. For general systems, the analogous target is dependence on the represented solution geometry instead of the full group $(G \oplus H)^d$.

Research question V.181.4 (Fixed size after nonuniform filtering). The prescribed-size theorem starts from uniform sampling on a fixed cardinality class. Conditioning a heterogeneous Bernoulli filter on its total size gives a nonuniform, product-weighted distribution. Can its conditional energy be maintained with a degree- $2s$ support state and no cardinality-sized auxiliary state? The present one-variable hypergeometric functional does not apply to unequal odds. A successful answer would combine the source’s strong phase filter with a hard size constraint without an unproved independence approximation.

Research question V.181.5 (Several genuine structural constraints). Can respected-relation surplus be retained together with prescribed Fourier correlations or energy on a specified spectrum? Linear combinations are easy to evaluate, but a positive combined objective need not make each constraint true. One needs a theorem producing a valid separating or pessimistic certificate with explicit margin and cost. Exact colour quotas may provide a useful intermediate model, but are not themselves a spectral-preservation theorem.

Research question V.181.6 (A sharp filter–precision–search frontier). The proof retains the slack $d = \alpha + \eta - \alpha\eta$ and uses a convenient common grid Q . Can exact influences, alternative positive kernels, and a smaller phase family be jointly optimized to improve the certified retained count at a fixed computational budget? The mathematical objective should specify both analytic loss and evaluation cost; optimizing only L inside one sufficient bound does not solve the full problem.

Research question V.181.7 (A faster exact array update). The branch probe uses only $O(s^2)$ entries, but maintaining every lower coefficient costs $O(s^4D)$ per label. Can batched translations, a Fourier representation, or a data structure for delayed updates reduce the total nD factor without destroying exactness? The target should include a bit-cost bound; numerical Fourier transforms with uncertified signs would not suffice for an exact monotonicity certificate.

Research question V.181.8 (Best possible use of a support profile). The initial coefficients c_j determine the uniform-cardinality expectation curve exactly. Which constraints on these coefficients are forced by graph energies, beyond nonnegativity for unsigned counts and the obvious support bounds? Sharp constraints could produce new threshold-free restriction theorems by locating a positive point on (II.191.5), rather than searching for a phase filter. The question concerns realizable additive profiles, not arbitrary polynomials of degree $2s$.

Research question V.181.9 (Proof-carrying implementations). Formalize the insertion and response invariants, and design a compact certificate format whose verifier is materially cheaper than rerunning the entire search. The final set and exact energy counts already certify quality; what remains is an efficiently checkable connection to the analytic retained-count guarantee and, where useful, to the searched maximum. A verified search need not be implemented inside the theorem prover if its output certificate has a small trusted checker.

Research question V.181.10 (Propagation beyond the local restriction step). Which later uses of Gowers’s restriction machinery can consume the constructed subset without requiring an infeasible enumeration of higher arrangements or a loss of required spectral information? The general-system theorem supplies a computational counting interface, not a replacement for those later structural estimates. A complete algorithmic version of the larger argument should track both quantitative density losses and execution costs through every stage.

V.182 Conclusion

Source 48, delivered Section 13.

The obstruction to direct conditional-expectation extraction is not the averaging principle. It is the interaction between repeated tuple entries and the structured statistic being averaged. A product of label-level exponential factors records each label exactly once, while divided powers retain the ordered multiplicities needed for additive energies.

The resulting state is stable under probability replacement. Universal Stirling response polynomials turn each replacement into weighted group translations, and a small zero-coordinate probe determines its exact slope. Recording distinct undecided support instead of total chosen size extends the same idea to fixed-cardinality sampling. Squarefree position variables then supply a general-system version, with repeated column types explaining the much smaller cost of balanced energies.

These identities give a fully explicit answer to the companion’s algorithmic extraction question in the stated ambient model. Together with certified rational filters and the retained analytic slack, they make the polynomial-loss restriction theorem constructive without a new loss in its retained-count bound. The more ambitious targets are now specific: phase search, sparse dimension dependence, nonuniform fixed-size filtering, simultaneous structural constraints, and formal proof-carrying execution.

V.183 Algorithmic skeletons and invariants

Source 48, delivered Appendix A.

V.183.1 Bernoulli signed rounding

The input is a shared label list, one image list and rational weight for each energy, an order, and rational probabilities. The following is mathematical pseudocode; the executable version is in `code/energy_inventory.py`.

```
Initialize one exact inventory for each image list.
J := the signed sum of its top, zero-group coefficients.
```


For each label i whose probability p_i is strictly between 0 and 1:
 Compute Δ_i from the universal response polynomials.
 Choose $q_i = 1$ when $\Delta_i \geq 0$, otherwise $q_i = 0$.
 Replace the i -th factor in every inventory, descending in degree.
 Assert $J_{\text{new}} = J + (q_i - p_i) \Delta_i$.
 Record the decision and continue with $J := J_{\text{new}}$.
 Return the labels with final probability one.

The central invariant is the complete coefficient identity (II.188.3), not only a numerical estimate for the top entry. Because every branch is an actual product distribution, the current state remains an exact conditional expectation.

V.183.2 Prescribed-size signed rounding

Initialize integer support-profile inventories with all labels undecided.
 Set $N := n$ and $r := \text{required cardinality}$.
 For each undecided label i :
 Compute each feasible exclude/include branch by the formal ratio.
 Evaluate exclusion using $L_{-}(N-1, r)$.
 Evaluate inclusion using $L_{-}(N-1, r-1)$.
 Choose a feasible branch of maximum signed value.
 Apply its profile update, descending in tuple degree.
 Decrease N , and decrease r exactly when inclusion was chosen.
 Return the included labels.

The invariants are the distinct-undecided-support interpretation, $0 \leq r \leq N$, and the exact two-branch averaging identity. The code tests that averaging identity at every nonforced branch in its regression instances.

V.183.3 A direct verification alternative

For small inputs, enumerate subsets and compute deterministic energy from sum multiplicities. If


$$r_{s,S}(h) = \#\{(i_1, \dots, i_s) \in S^s : g_{i_1} + \dots + g_{i_s} = h\},$$

then $E_s(S; g) = \sum_h r_{s,S}(h)^2$. This formula gives an independent reference for the final set and, by weighting every subset with its exact Bernoulli probability, for the initial expectation. For partial coefficients, the test suite also lists ordered tuples directly and multiplies probabilities once per distinct support. Neither reference uses the response identity being tested.

V.184 Proof-dependency map and audit boundaries


Source 48, delivered Appendix B.

The purely finite algebraic part has the dependency chain



The prescribed-size extension adds distinct-support counting and a hypergeometric functional; it does not use Fourier analysis. The general-system theorem changes only the coefficient algebra and its index set.

The source theorem adds a separate analytic chain:



This separation makes the boundaries explicit. A future better filter can reuse the same exact rounding theorem. A future faster array implementation can preserve the same analytic statement. Neither change by itself supplies the independent structural ingredients needed to improve a global Szemerédi bound.

V.185 Mathematical consequences and integration boundaries

Source 49, delivered Section 9.

V.185.1 What the degeneracy improvements change

The prior two-cube argument treated every nonaffine certificate as a possible codimension-two locus. Theorem II.222.2 shows that this worst case contributes no new locus at all: its translations are already in the affine divisor. This is why a new asymptotic coefficient can be determined without enumerating every high-degree ternary polynomial. The principle is useful more generally. A count of certificates by rank alone can overestimate the union substantially; one should ask whether the low-rank loci are geometrically contained in components already accounted for.

The exact $k = 3$ calculation adds a second principle. A prime-by-prime search can be replaced by an integer obstruction certificate. Once the top coefficients determine a candidate translation, a gcd of finitely many integer residuals identifies every possible exceptional characteristic at once. Small-prime symmetry then finishes a finite list. The fact that the surviving cubic gcds have only primes two and three, and the quadratic ones only three and seven, is stronger evidence than many successful large-prime trials; it is the reason the result is a theorem about every odd prime.

Neither principle automatically applies to the multiple-slice case. For $d \geq 2$ the shared last-coordinate equation permits more patterns and the geometry of its vertical moments differs. A future claim about $\delta_{k,d}$ must start from those equations. Similarly, replacing a local degeneracy probability in a Riesz-selection argument requires retaining the weights and threshold conditions in that argument. The present exact two-cube identities are not silently substituted for the eight-arrangement interface I3 or I4.

V.185.2 What the norm improvement does and does not imply

The odd-order and mean-zero hypotheses have specific roles. Odd order makes a nonzero real Fourier coefficient pair with a distinct inverse coefficient. Centering removes the trivial character and constant harmonics of the dual test. If these assumptions are dropped, the constant function or a real order-two character has all the relevant uniformity norms equal, so a factor greater than one is impossible without a replacement hypothesis.

The stronger non-five statement is pointwise: it is enough that one largest Fourier coefficient has order different from five. The group may still have five-torsion. Conversely, the weaker order-five dual certificate does not show that five-torsion truly lowers the optimal norm factor. It only locates the loss in the particular two-harmonic test used here. Extra information about the remaining Fourier spectrum could improve its tail bound.

The Fourier concentration envelope in (III.93.20) contains more information than (III.93.4). It can be used whenever an application independently constrains the largest Fourier pair. The strict inequalities for individual nonzero functions do not imply strict inequalities for the supremum constants: a sequence of groups and functions could approach a limiting extremum. No compactness or attainment theorem is assumed.

There is broader work on uniformity norms, including recent estimates for their L^p variants [71]. Our target is the precise real, centered comparison on finite odd-order groups and the explicit constants relative to the project's baseline. A comprehensive priority comparison or a classification of extremizers would be an additional research task. The algebraic proof and the stated constant do not depend on such a priority assertion.

V.185.3 Why the model theorem has two dimensions and four controls

In the cover theorem, active coordinates specify the variables on which a model depends. Ambient coordinates specify the box in which the partition must live. A partition in the active coordinates may have a common difference incompatible with the untouched sides of an ambient box. Keeping the ambient dimension in the induction resolves this directly. It also prevents an exceptional set from being assumed to be cylindrical when the proof has not made it so.

The four controls describe different mathematical requirements. N_j is a field-size condition needed to choose a global core. Q_j bounds the number of models on a local cell. $c_{j,D}$ is a power in the resulting width, and $T_{j,D}$ says when that power guarantee starts. Combining them into one number can hide a dependency on the local error inside a supposedly global choice. Their explicit separation is what permits the quantifier order in theorem II.235.5.

The improvement in the local-error exponent table comes from retaining the actual dimensions of the proper-section terms. A cube remainder contains $\binom{k}{l-1}$ terms of active dimension l , rather than $2^k - 1$ terms all of dimension k . The resulting better exponents are exact consequences of this recurrence. Their rapid growth also shows that replacing successive refinements by a genuinely simultaneous theorem remains a substantial problem.

V.186 Further research questions and a proposed program

Source 49, delivered Section 10.

The questions below are left unresolved by the present manuscript. They are proposed research targets, not claims that each formulation is a recognized open problem in the published literature. The first group continues the earlier project's exactness question; the second continues its real norm extremal problem; the third concerns the new quantitative recurrence and its integration.

V.186.1 Translation geometry and higher degeneracy coefficients

Question V.186.1 (Exactness beyond dimension three). *For $k \geq 4$, is $T_k(p) = \mathcal{B}_k(p)$ for all sufficiently large primes? If not, what is the smallest dimension, characteristic, and derivative rank admitting a translation outside the signed-coordinate arrangement?*

The immediate milestone is dimension four. A rational off-divisor certificate that persists modulo all sufficiently large primes would settle the eventual exactness conjecture negatively; an isolated

exceptional-prime example would refute only exactness at every prime. A positive proof should classify or compress the low-rank translation loci rather than compare all pairs of 3^{16} patterns. A useful certificate for a counterexample consists of the two ternary vertex tables, the translation, and all squarefree coefficient identities.

Question V.186.2 (The exact cubic-order coefficient). *For fixed k , determine the coefficient of p^{-3} in $\delta_{k,1}(p)$. Which codimension-three translation loci actually contribute outside $\mathcal{B}_k$, and with what overlap multiplicities?*

Theorem II.224.2 ensures that a fixed coefficient exists for large primes. The first two coefficients required only containment of rank two. The next coefficient tests rank three and intersections with the known divisor. Counting patterns alone will not distinguish coincident loci or loci already contained in the divisor.

Question V.186.3 (Derivative-rank-three classification). *Determine the optimal number of essential coordinates for a ternary-valued multiaffine polynomial of derivative rank three, with explicit characteristic restrictions, and classify the extremal cases.*

The rank-two answer is four, and its exact census is a polynomial in k . A comparable classification at rank three would make the third-coefficient problem much smaller. Degree-three constraints alone do not retain all the linear dependence among the partial derivatives; the interaction of that dependence with three-valued Boolean images is the relevant structural information.

Question V.186.4 (A sharper finite remainder). *Replace $M_k^2/2$ in (II.222.6) by an explicit count of distinct off-divisor affine loci, or by an incidence bound with a substantially smaller dependence on k .*

Common negation is only the first available symmetry. Boolean coordinate reflections, permutations, and repeated translation orbits may identify many more pairs. The desired statement should bound the geometric union, not merely improve a census of certificates that may define the same subspace.

Question V.186.5 (Exceptional characteristics and characteristic polynomials). *Compute χ_4 or a nontrivial initial segment of its coefficients, and obtain useful bounds on the bad-prime set $\mathcal{E}_k$ in theorem II.224.2.*

The generic elimination algorithm is effective but too large for naive use. A better method would keep only translation strata that can affect the desired coefficients and use rank bounds to discard the rest. The cubic gcd certificate suggests that arithmetic obstructions may be much smaller than general determinant estimates predict.

Question V.186.6 (Human proofs of the three-dimensional certificate). *Find a short classification explaining the residual gcd histograms in lemmas II.223.2 and II.223.3, without a full enumeration of all vertex patterns.*

This is a proof-compression problem rather than a missing step in the finite proof. An informative answer would explain the special primes three and seven and the scarcity of rational off-divisor candidates, and would suggest which part of the argument should extend to dimension four.

Question V.186.7 (The multiple-slice second coefficient). *For $d \geq 2$, determine the second coefficient of the arrangement degeneracy probability while retaining both horizontal and vertical moment constraints.*

One should first classify the codimension-two loci that are not already forced by simpler coordinate coincidences. The present two-cube result shows why a broad union bound can miss an exact coefficient, but its vertical collapse is special and cannot supply the answer by itself.

V.186.2 Dual tests and sharp real norm comparison

Question V.186.8 (The fourth-order extremal constant). *Determine c_4^{odd} , or reduce the interval in (F.54.7) substantially. Does the supremum arise as a limit of functions on cyclic groups with a bounded number of Fourier pairs?*

An optimization over finitely many harmonics produces candidates, but a sharp result also requires a dual inequality for every spectrum. The present proof illustrates both sides: a test inspired by a lower example yields a universal upper certificate, but its tail estimate still loses information.

Question V.186.9 (The role of five-torsion). *Is the optimal fourth-order constant on groups with no five-torsion different from the unrestricted odd-order constant? Can the order-five dual profile be raised to the non-five level?*

The current gap between $29/4$ and $15/2$ is a gap between available certificates. It does not prove different extremal constants. A useful first step is an exact finite-dimensional optimization on the five-point quotient that retains the phases of both nontrivial inverse pairs.

Question V.186.10 (A family of optimized dual functions). *For $g = \phi_1 + t_3\phi_3 + t_5\phi_5 + \cdots + t_{2s+1}\phi_{2s+1}$, find rational coefficients and a tractable exact certificate that improves the principal-to-tail ratio of D_4g .*

The challenge is to improve the test's dual tail and its own norm simultaneously. A numerical search should be followed by rational interval certification and complete treatment of the finite set of small character orders at which step-frequency constraints alias. The finite formula in (III.93.11) provides a precise starting point.

Question V.186.11 (Beyond the largest-pair parameter). *Replace the single mass p by a profile involving the two or three largest inverse pairs. Can their relative orders or additive relations force a stronger moment-dual lower envelope?*

The current dual tail bounds each unwanted coefficient separately by the total remaining fourth mass. Equality in all those bounds would usually demand incompatible concentration patterns. Keeping several pair masses could measure that incompatibility, but modular aliases must still be included.

Question V.186.12 (Higher orders and extremizing sequences). *Develop direct dual tests for U^d at $d \geq 5$, and determine how c_d^{odd} and near-extremizing Fourier profiles behave as d grows.*

The U^5 improvement here combines a fourth dual test with a higher moment. It leaves room for a genuinely fifth-order dual certificate. Any asymptotic conclusion should distinguish a supremum over all groups from behavior on a fixed group and should account for the possible nonattainment of a supremum.

V.186.3 Model-theorem controls and formal integration

Question V.186.13 (Simultaneous anchor refinement). *Can all anchor slices be modeled together with a width loss that grows polynomially, rather than exponentially, in the number of anchors?*

The factor c_H^r in (II.235.28) is the precise place to improve. A theorem must preserve the global core's independence of the local tolerance and work in the same ambient common-difference box geometry. Improving the sampling count alone does not resolve this refinement cost.

Question V.186.14 (Sharper phase partition input). *Replace I7 by the strongest suitable simultaneous multilinear recurrence theorem and recompute the exact Q, c, T profiles. Which improvement in I7 changes the final width exponent rather than only the threshold?*

Because the induction is written as an explicit implication, a stronger interface can be substituted without rewriting the global extraction. The graph count is independent of the spectral phase-model count, whereas the width and threshold are not. Tracking that distinction is necessary to state a real quantitative gain.

Question V.186.15 (Propagation through the remaining density increment). *Assuming the seven interfaces have been independently certified, carry the new Section 16 controls through the later phase-extraction and density-increment arguments and state the resulting complete bound.*

This is the step needed before claiming any new global Szemerédi estimate. The first deliverable should be an explicit recurrence with all positivity, minimum-width, and prime-size hypotheses, even if its growth is weaker than a desired closed form. A successful simplification must then be proved from that recurrence.

Question V.186.16 (Formal certificates and a modular theorem ledger). *Formalize the translation reduction and rank-two classification, then verify the two finite certificates in a proof assistant and represent the cover recurrence by a dimension-indexed expression type.*

The natural order is structural lemmas first, finite integer reflection second, and the deep imported interfaces as explicitly named hypotheses. Termination of the new recurrence is a simple decrease of active dimension. Proving it does not require hiding a failed numerical domination inside a general automation tactic.

V.187 Reproducibility and repository integration

Source 49, delivered Section 11.

V.187.1 The exact verifiers

The package contains three independent exact entry points, plus an optional numerical cross-check. The exact programs use only Python’s standard library and reject optimized execution that would disable assertions. They regenerate their JSON output beside the scripts. No network access, external computer algebra system, or random seed is required.

Program	Mathematical task
<code>verify_rank_two.py</code>	Regenerates the square-invariant table; checks derivative ranks and the census in small prime fields; verifies exact two-dimensional counts, three-dimensional rank-two translations, and the displayed first two coefficient formulas in representative dimensions. The all-dimension theorem is proved in the text, not inferred from these supplemental tests.
<code>verify_dimension_three.py</code>	Performs the complete cubic and quadratic integer obstruction enumerations; checks their formulas by direct translated-polynomial expansion; exhausts all fourteen signed-permutation representatives at primes eleven and thirteen. Together with the text’s reduction, this is the finite proof of theorem II.223.1.
<code>verify_dual_gap.py</code>	Regenerates every fourth-dual coefficient polynomial for orders 5, 7, 9, 11, ∞ ; forms all phase-dependent norm and alias bounds; checks every rational table bound and every comparison used in the fourth- and fifth-order conclusions; regenerates the order-three full-cube count.
<code>check_dual_numeric.py</code>	Optional independent direct average of punctured-cube products at five character orders, two phases, and three basepoints. Uses NumPy and floating point, and supplies a cross-check only. It is not a premise of a theorem.

The dimension-three program examines 1 985 589 ordered cubic pairs and 12 610 ordered quadratic pairs with the same nonzero quadratic part. Its small-prime part tests $14 \cdot 6561 = 91\,854$ polynomial/representative cases. The dual verifier uses five face-split expansions, each with $4^8 + 4^7$ face assignments, and a separate 2^{16} sign enumeration for order three. These are finite exhaustive calculations at the reductions where the article explicitly invokes computation. They do not establish the imported interfaces I1–I7.

V.187.2 Build and file layout

The top-level `gowers_rank_dual_refinements.tex` is a single complete source file, including its bibliography and the vector plot. It does not depend on the earlier project manuscripts being present on the build machine. The bibliography identifies those dependencies mathematically. Compile with a standard TeX Live installation including `pgfplots`; the package’s `Makefile` invokes `latexmk -pdf`. Running `make verify` regenerates the three exact certificate files. Running `make numeric` is optional.

The proposed repository destination is a new subdirectory of the Ramsey research directory. It should be integrated as a research manuscript with computational certificates. Its conditional cover theorem should retain the I1–I7 hypothesis list. A theorem ledger should not mark the article’s finite proofs as Lean-certified before an actual formal implementation is checked. Existing source labels and earlier manuscript numbering can remain intact; all new internal labels here use descriptive prefixes.

V.187.3 Source pin and limitations of the audit

The repository was inspected at commit `8c12c906e4ab5a00001a65de3c882f935a15661a`. The source-declaration locations for the structural interfaces are recorded in table II.235.1. They are locations in that transcription, not proof-status assertions. The earlier project reports cited below were also read to avoid rediscovering their sampling, interpolation, moment, and first-order degeneracy results.

The mathematical development was independently reviewed across the rank, duality, small-characteristic, and recurrence arguments, and the exact programs were rerun when assembling the deliverable. These checks support the manuscript’s stated proof status. They are not a substitute for peer review, and they do not make an unproved deep interface into a formal theorem. No repository files were changed as part of preparing this package.

V.188 Computational checks and reproducibility

Source 50, delivered Section 10.

The proof of Theorem II.233.1 is mathematical; no computer can realistically enumerate its displayed field. The accompanying programs serve two narrower purposes: exhaustive verification of the finite cover inequalities at small moduli, and a reproducible finite illustration of the same obstruction. They use exact integer and rational arithmetic and require only Python’s standard library.

V.188.1 Exhaustive checks of the finite inequalities

The script `code/verify.py` checks all $3^5 = 243$ words over a three-symbol alphabet in $\mathbb{F}_5$. There are 31 distinct nonempty modular progression sets in this field. For every such word and set it checks all affine lists of length zero, one, or two, allowing repetitions in the two-element lists. It also checks all single quadratic polynomials. A separate test enumerates every two-symbol word on $\mathbb{F}_3$, every product of nonempty progression sets, and every pair of bilinear polynomials. A further quadratic test uses all $2^7 = 128$ two-symbol words on $\mathbb{F}_7$ and its 99 distinct progression sets. Finally, all 3^9 maps on $\mathbb{F}_3^2$ are tested for equivalence of separate affinity and bilinearity. These are exhaustive finite checks, not statistical samples.

Test class	Exact checks passed
Affine lists on $\mathbb{F}_5$	2,644,083
Single quadratic graphs on $\mathbb{F}_5$	941,625
Bilinear graph pairs on product boxes over $\mathbb{F}_3$	1,301,832
Single quadratic graphs on $\mathbb{F}_7$	4,346,496
Separate-affinity characterization on $\mathbb{F}_3^2$	19,683
Reduced-exponent integer budget comparisons	54

Some of the root envelopes on short sets are vacuous. In the $\mathbb{F}_7$ quadratic test, 1,205,302 checks have an upper bound strictly below the size of the tested set. The endpoint test finds exactly 81 separately affine functions, all bilinear; it does not test every real-weight instance of the product property.

The reduced-exponent tests do not instantiate $E = 2^{1024}$. They test the algebraic budget comparison at small exponents; the all-exponent statement is proved in Section II.233. Passing these tests is not represented as a formal proof of the manuscript.

V.188.2 A complete small branch-obstruction certificate

The same script produces a specific four-symbol word on $\mathbb{F}_{257}$ using the recorded pseudorandom seed 20261006, then checks *every* modular progression parameterization of length at least 128. The resulting exact extrema are

$$\max_{I,j} \frac{b_j(I)}{|I|} = \frac{27}{64}, \quad \min_{I,j} \frac{b_j(I)}{|I|} = \frac{11}{138}. \quad (\text{V.188.1})$$

This involved 8,552,960 progression parameterizations and 34,211,840 symbol-frequency checks. The maximum occurs, for example, for symbol 0 on the length-128 progression with start 135 and difference 19.

Every symbol has at least four occurrences on every checked progression, as verified directly in the script. Theorem II.228.2 therefore identifies the best affine lists of each length $q \leq 4$ with the q heaviest constants. In particular, a single affine graph covers at most $27/64 < 1/2$ on every such progression. By the product-box theorem, the function $F(h, y) = a(y)$ cannot be covered on half of a union of boxes with those final-axis lengths by one multiaffine graph per cell, even though four constant graphs cover it exactly.

The word, all parameters, the extrema, and an extremal progression are stored in `data/verification.json`. This is a concrete small analogue of the branch-budget obstruction. It is **not** the enormous instance of the repository's particular control functions.

V.188.3 An exact deterministic construction test

The script `code/construct_word.py` implements conditional expectation for a uniform alphabet, representing all conditional binomial tails by integers on a common denominator. Its default run uses

$$p = 31, \quad R = 2, \quad L = 20, \quad \varepsilon = 2/5.$$

It checks 4,682 distinct progression sets. The initial expected number of bad progression-symbol events is exactly

$$\frac{65,486,981}{134,217,728} < 1.$$

The final number is zero, and the maximum discrepancy of the output word is exactly $2/5$. Here the exact initial expectation is better than the general union-bound certificate in Theorem II.230.2; the algorithm needs only that the actual expectation be less than one. Its output is recorded in `data/construction.json`.

V.188.4 Commands and deliverables

Run the checks from the package root with

```
python3 code/verify.py
python3 code/construct_word.py
make pdf
```

The package includes the manuscript source and compiled PDF, both programs, the exact output files, a source-audit document, a formalization plan, and a Makefile. No unproved Lean file is included. The optional arguments of the programs are documented by `-help`. Large moduli can make the exhaustive checks expensive; the construction algorithm's polynomial dependence is on p , not on the bit length of p .

V.189 A formalization route

Source 50, delivered Section 11.

The target is a theorem refuting the pinned catalogue entry, not a patch that silently weakens its conclusion. The cleanest development would separate a small finite combinatorial library from the source-specific witness assembly.

V.189.1 Finite combinatorial core

Formalize the root bound for a nonconstant polynomial agreeing with an R -valued function on a finite set. Then formalize the finite list envelope using a union bound and the maximum mass of an s -element subset of the alphabet. The latter formulation avoids sorting until a separate equivalence with order statistics is useful. Exact-optimum and deficit lemmas follow from one inequality comparing the s - and q -branch masses.

Next prove the fibrewise product estimate by summing over the base set and the adaptive partition estimate by summing over disjoint carriers. These steps do not require Gowers norms, Fourier analysis, or the inverse theorem. The counterexample uses only these finite counting statements after the word has been constructed.

V.189.2 Existence of balanced words

A formal proof can use independent coordinates on the finite probability space $S^{\mathbb{F}_p}$ with uniform measure. Only Bernoulli concentration and a union bound over at most Rp^3 events are needed. Alternatively, define the exact bad-event score used by `construct_word.py`; a finite averaging argument shows that some word has score below one once the probabilistic estimate is established. There is no need to formalize an asymptotic random process or uncountable probability space.

The main theorem uses a rational tolerance $1/(4R)$, a natural integer R , and a finite prime field. Real logarithms occur only in a sufficient size threshold. One can first formalize a conditional theorem taking a balanced word as input, then add the finite existence theorem, and only afterwards formalize the convenient explicit prime threshold.

V.189.3 Source-specific assembly

The witness assembly has four parts:

- (a) Build the full proper two-dimensional box and show its width and cardinality are p and p^2 .
- (b) Set the cube-selection family to the finite universe and prove that `section16GoodDomain` is the universe and `section16PhiOne` is $(h, y) \mapsto a(y)$.
- (c) Prove the identity-partition lemmas for constant and finite-range functions, including width zero, and use them to satisfy both premises.
- (d) Specialize the conclusion to loss $1/2$, derive the Q -graph and $p^{1/Q}$ -width requirements, and contradict the retained mass by summing cellwise covers.

The completed statement could then be expressed as the negation of the pinned `lemma_16_10`. Until those steps are kernel checked, that is a proposed formal theorem, not a delivered formalization. The manuscript's mathematical proof and the finite test suite are evidence of different kinds and must not be conflated.

[write] **7 October 2026; sources 47–53.** *Done in the repository* (checked at HEAD 06b30ca52): the negative theorem is formalized as `lemma_16_10_packaged_premises_counterexample` (note in Section F.55), by the four parts (a)–(d) above, with the finite core in `Proofs16FiniteAlphabetBoxCover.lean` and the balanced words in `Proofs16BalancedFieldWord.lean`; the catalogue entry is preserved and stays open, as this section recommends.

V.190 Further research questions

Source 50, delivered Section 12.

The following questions are not assertions of established open status in the wider literature. They are specific next problems arising from the proved obstruction and the audited repository interface.

Research question V.190.1 (An upstream hypothesis that genuinely excludes the example).

Proposition II.234.1 settles the full-domain endpoint $\gamma = 1$: the product property excludes the witnesses. For $0 < \gamma < 1$, identify the quantitative information retained by the relation product property after selections and restrictions. Does it adequately control one-coordinate summands at the branch budget needed for the induction? A useful answer must give explicit bounds; referring generally to “more structure” is insufficient. Adding both orientations of one-dimensional section control alone does not work by Proposition II.233.3.

Research question V.190.2 (A sound replacement for the induction control). Starting from the explicit local lift, derive a closed recurrence for graph counts, width exponents, and iteration parameters that remains valid after all losses and common refinements. What is the smallest additional dependence on the input branch count forced by that recurrence? The uniform examples require at least $(1 - \eta - o(1))R$ graphs to retain $1 - \eta$ of long boxes, so any proposed compression should first pass this lower-bound test.

Research question V.190.3 (Branch profiles instead of a single worst-case count). Can the interface carry the whole function $q \mapsto M_q$ rather than a single branch count? For nonuniform examples, discarding many low-mass branches can be inexpensive, while a uniform profile cannot be compressed. Develop composition rules for such profiles under union, interpolation, and common partition refinement, with an explicit treatment of correlations between the branches on different fibres.

Research question V.190.4 (The shortest universal balance scale). For fixed R and prescribed relative error, determine sharp constants or orders for the least length L at which some word is balanced on every modular progression of length at least L . Our absolute-balance construction uses a logarithmic scale for fixed parameters, but no optimal constant is claimed. A lower-bound analysis should distinguish modular progressions from ordinary intervals and distinguish a fixed alphabet from a growing alphabet.

Research question V.190.5 (Optimal scales for polynomial-list rigidity). Theorem II.231.4 gives a sufficient scale $O(R(d+1)\log p)$ at fixed relative error. Determine the actual transition at which all optimal degree- d lists become constant lists. The sharp implicit nonconstant-graph threshold is smaller than our convenient T when d is fixed, but relative symbol balance may dominate. Which of these two phenomena sets the true threshold in each growth regime for R and d ?

Research question V.190.6 (Beyond the exact-optimum regime). For a fixed finite word, classify the list-cover optimum when $b_{(q)} < dR$. The deterministic envelope need not be attained there. Can one replace the root allowance by a computable incidence invariant that gives an exact formula or a certified approximation without enumerating all lists? Even the affine case could reveal a useful finite optimization problem.

Research question V.190.7 (Efficient explicit constructions). The conditional-expectation algorithm is polynomial in p but not optimized. Can the simultaneous progression constraints be certified and updated in near-quadratic or near-linear time in p for fixed alphabet? For the stronger polynomial-avoidance construction, can the coefficient-space union bound be implemented implicitly rather than by enumerating p^{d+1} polynomials? An explicit algebraic word with logarithmic-scale balance would be particularly useful for formal counterexample libraries.

Research question V.190.8 (Torsion and non-field codomains). Over composite rings a nonconstant affine map may have large fibres, and the root-count proof changes. Determine the exact replacement of the cover law in terms of fibre multiplicities or annihilator sizes. Can one identify a profile invariant that behaves well under products of cyclic groups and reduces to M_q over prime fields?

Research question V.190.9 (Stability from incomplete observations). The global stability theorem assumes a known alphabet and exact values of the word. Extend it to noisy observations, partially observed fibres, or probability-weighted point sets. The Hamming robustness corollary handles adversarial edits after construction, but it does not yet give a reconstruction algorithm for the underlying branch profile.

Research question V.190.10 (A certified catalogue correction). Formalize the negative theorem and decide how the catalogue should distinguish the printed contextual statement, the false two-input abstraction, and a replacement local lifting theorem. The source audit should preserve the old statement for traceability while preventing it from being counted as an available lemma in the main theorem’s dependency graph.

V.191 Conclusion

Source 50, delivered Section 13.

The exact finite cover law identifies the obstruction without relying on numerical experimentation: a sufficiently populated constant branch cannot be replaced, at equal coverage, by a low-degree nonconstant polynomial. Balanced words make this true simultaneously on every long progression, so adaptive partitions cannot eliminate the obstruction. The profile realization and growing-alphabet theorems quantify both the mass frontier and the scales on which it is visible.

Applied to the pinned Section 16 definitions, this yields an actual counterexample to the two-input form of Lemma 16.10, at fixed parameters and arbitrarily large prime moduli. Every named witness and every input budget has been checked in the manuscript. The result demands a revision of that abstraction, not a reinterpretation of a failed calculation as an unproved but harmless estimate. Positive lifting remains available when its genuine branch and width parameters are retained.

The outstanding larger task is to determine which additional structure in the full proof can justify a suitable replacement, and at what quantitative cost. That task is sharply separated here from the local statement that has been resolved.

V.192 Source-to-witness crosswalk

Source 50, delivered Appendix A.

Pinned definition or input	Witness or consequence in this article
N, k, theta, gamma	$p, 1, 1, 1$, with p prime beyond (II.233.2).
B, H1, x0	Full $\mathbb{F}_p^2$, full $\mathbb{F}_p$, and 0.
Y h	The complete finite type of cube elements with sidelength h .
Section16GoodInducedPair	The horizontal cube based at zero exists for every (h, y) and belongs to Y_h .
section16GoodDomain	The full set $\mathbb{F}_p^2$.
section16PhiOne	$\phi(0 + h, y) = a(y)$.
crossSectionR	$r_0 = 16^{2^{128}} \geq 2$.
FinalCoordinateSections	Every section is a constant, certified by the identity partition.
MultiplyLinear	$(2r_0/\sigma)^{Er_0} \geq R$ for every $0 < \sigma \leq 1$.
section16Lemma9QBound	At least 1; choose $q_\delta = 1$.
section16Lemma9DeltaQBound	$(\zeta/2)m^e \leq m$ for $m \geq 1$, with $0 < e \leq 1$; zero for $m = 0$.
section16Lemma9Width	Identity partition, retained set P , R constant affine maps.
Section16AllBoxLineCovers	At loss $1/2$, it demands at most Q graphs and widths at least $p^{1/Q}$.
MultiplyLinearFunction gamma 1	

Pinned definition or input	Witness or consequence in this article
Target retained set	Must contain at least $p^2/2$ points, but every allowed union covers at most $5p^2/16$.

V.193 Definition-level pseudocode for the witness

Source 50, delivered Appendix B.

The following is mathematical pseudocode, not a claimed Lean proof. It exposes the finite objects and the universal input witnesses.

```

E := 2^(2^10)
Q := 2^E
R := 4*Q
c := 1/Q
p := any prime greater than 2^(4096*Q^4)
S := {0, ..., R-1} as distinct elements of F_p
epsilon := 1/(4*R)
L := ceil(8*R^2*log(4*R*p^3))
a := an AP-balanced word supplied by conditional expectation

B := all of F_p^2
H1 := all of F_p
x0 := 0
Y(h) := every cube element with sidelength h in B
phi(h,y) := a(y)

For every fixed y and every input box:
  retain the entire box
  use its identity partition
  use the single constant graph a(y)

For every sigma, m, and proper input box P:
  qGamma := R
  qDelta := 1
  retain P
  use the identity partition P
  use final-coordinate affine maps y |-> v, for v in S

To refute the target:
  take P := F_p^2 and rho := 1/2
  every allowed cell has final-axis length at least p^c
  every allowed list has at most Q affine restrictions
  total covered fraction <= Q*(1/R + epsilon) = 5/16
  required retained fraction >= 1/2

```

V.194 Proof and provenance checklist

Source 50, delivered Appendix C.

The source-dependent statement has been pinned to an exact commit and blob. The named domain B_1 has been realized using actual cube witnesses. The constant-function certificates cover all proper input boxes, all allowed loss parameters, and the zero-width edge case. The two line-cover count parameters are kept distinct. The conclusion is tested on a full proper box, and its integer graph count is bounded by the source’s real control without changing the inequality direction. Both the balance scale and the root-counting scale are dominated by the permitted output width. The pointwise cover is summed only over disjoint cells. The Hamming robustness claim preserves the alphabet and therefore preserves the two input certificates.

The positive local-lift theorem is attributed to the source, not to this article. The finite verification scripts are explicitly separated from a kernel proof. The new manuscript does not assert that the selected relation satisfies every upstream product-property hypothesis, does not claim an error in Szemerédi’s theorem, and does not advertise a new global progression-free bound. The further questions are research directions rather than a claim that all of them have been certified open in the wider literature.

V.195 The two interfaces after complete upstream extraction

Source 51, delivered Section 11.

The two results have different quantifiers. Theorem F.56.1 proves that even the full budgets and all retained predicates do not force a grid for arbitrary selected outputs. Corollary II.211.4 constructs suitable outputs by preserving information in the actual selection, then localizing before the final square extraction. In the zero-phase counterexample other bilinear squares were already easy to choose; the constructive result works with the original phase and domain of a general admissible context, using the stated analytic source theorem.

The latest source already proves complete square extraction with a density-uniform threshold as described in Section II.210.6. The following items describe a formal interface for the particular direct route proved here, not missing prerequisites for every current upstream square construction:

1. Export the actual fact $\ell(R) = m$ or $m + 1$, or just $\ell(R) \leq m + 1$, through the existing Stage 13.6 wrappers. It belongs to their selected witness, not to every object satisfying the old conclusion.
2. Record short-carrier rectification with optional reversal and unchanged carrier, and the balanced aligned-parent partition with lengths $\lceil L/2 \rceil$ through L .
3. Reuse the existing all-scale coefficient localization at unchanged density, rectify its selected U inside the short parent, and feed the recovered stride and side $\ell(U) - 1$ into the common-step cover.
4. State the direct Stage 13.8-to-square result with its effective width (II.211.9), rather than claiming the old Stage 13.9 lower length or the arbitrary-witness catalogue Corollary 13.10.

This report establishes the mathematical geometry and source-interface assembly; it does not claim those strengthened wrappers or the direct theorem have been newly implemented or Lean-checked. The existing context-to-extraction, covering, parent-localization, complete fixed-density square, and uniform actual-density square results are credited at their pins. Reconciliation with the unchanged catalogue or printed target statements, optimization beyond the stated effective bounds, and a new global quantitative Szemerédi bound remain outside the result.

V.196 Proof and computation have different roles

Source 51, delivered Section 12.

The proof of Theorem F.56.1 is symbolic; the constructive geometry is also proved for arbitrary admissible parameters, with its analytic input explicitly imported in Section II.211. Existence of one sufficiently large prime follows from the finite list of fixed positive-exponent inequalities; no prime at that scale has been computed. In particular, the package must not be interpreted as a numerical witness satisfying the six full budgets.

The Python companion uses bounded exact integer arithmetic to check finite exponent identities, small difference-set obstructions, endpoint and residue-chain covers, density averaging on bounded cases, convolution counts and fixed-height normalizations for small arrangement examples, and exact Fourier normalizations. Further bounded checks cover short-carrier rectification, balanced aligned-parent partitions, weighted density preservation, the side- $M - 1$ fit, and original-coordinate phase reversal. It also verifies the frozen source inventory against byte lengths, SHA-256 hashes, and Git blob identities. It does not approximate the astronomical exponents by floating point, expand $2^{9437184}$, build a denominator 2^T , search for the huge modulus, or enumerate the enormous graph family.

Finite checks are useful for catching an indexing error or a missing factor of N ; they do not prove the universal statements. Conversely, a successful source hash check verifies the identity of inspected bytes, not the logical correctness of a Lean theorem or the health of a whole project. No repository file was edited, no new Lean proof was produced, and no build or axiom audit of the upstream project is claimed.

The deliverable includes the editable L^AT_EX source, PDF, bounded companion, tests, frozen relevant Lean source snapshots and curated hash records, and an offline reproducible packaging command. The README describes the exact command-line limits, provenance scope, and output checks. A checksum inventory detects accidental changes; it is not an independent trust anchor for an archive whose bytes an adversary controls. The article supplies the mathematical proof at full scale, while the companion supplies deliberately bounded diagnostics.

V.197 Pinned source map and comparison scope

Source 51, delivered Appendix A.

All paths below are relative to `Combinatorics/Ramsey/Lean/GowersSzemerédi/` in [VladimirReshetnikov/ProveIt](#). The primary definition and budget pin is

`aad4789fcab0cc1d045d3247988596abfd4f5c2c`

The newer conditional-cover pin is

`330102ae0f0128b5569033a63166278f56179e3b`

The latter commit was observed in the frozen source check at 00:25 UTC on 7 October 2026. This is a dated comparison, not a claim about an unpinned moving branch after that observation.

Definitions.lean

[Lines 43–47](#) define an exact finite partition; [lines 81–98](#) define modular progressions, their carriers, and properness; [lines 135–137](#) fix the unnormalized Fourier transform. Centered absolute value and the Freiman/additive-tuple conventions occur at lines 278–279 and 311–321.

Section10.lean

Lines 434–455 define λ_{10} , the real spectrum bound, its integer ceiling, and the corrected ζ_{10} with division by that integer.

Sections12_13.lean

Lines 57–103 define arrangements and fixed-height counts; lines 198–259 define vertical edges, separate Freiman behavior, respect, and bilinearity. The complete context and stages are at lines 302–367, 397–407, 438–454, 479–500, 536–547, and 572–582. The unrestricted catalogue conclusion is at lines 594–607. These numbered catalogue claims are `Prop`-valued definitions; their presence alone does not assert a proof.

Proofs13BilinearExtraction.lean

Lines 16–34 contain the exact six budgets, initial-stage assumptions, and existential retained output interface. The proof ends at line 49.

Proofs13BilinearRestriction.lean

Lines 29–43 define `Corollary1310Grid`. Lines 111–139 derive product support from the actual stage chain, and lines 267–288 give the grid-qualified wrapper. The main theorem here obstructs its grid input, not its conditional proof.

Proofs13LargeScaleRecurrence.lean

Lines 120–150 contain the existing fixed-density context-to-extraction theorem. We do not apply it to the moving density in this report.

Proofs13CommonStepCover.lean at the newer pin

Lines 29–105 give the padded cover and capacity averaging; lines 107–150 give the general integer-fit bilinear-square theorem; lines 174–214 give its square-root Corollary 13.10 wrapper.

Proofs13ParentLocalization.lean at the aligned-parent pin

Lines 11–27 construct the short-parent partition; lines 30–90 contain the parent-size-qualified bilinear-square theorem summarized in Section II.210.3. Its source is frozen with a separate commit pin from the primary counterexample inputs.

The imported coefficient API and exponent definition

`Proofs13CoefficientPartition.lean`, lines 84–136, contains the all-scale coefficient localization and its Stage 13.9 exponent identity. The preceding lines 64–80 give the original-coordinate bilinear formula. `Proofs07SimultaneousLinearity.lean`, lines 1075–1077, defines the one-domain exponent. The all-scale coefficient theorem is imported as declared; the report does not reprove the analytic partition theory.

The actual short-parent witness

`Proofs13CommonStepSelection.lean`, lines 92–182, supplies the m or $m + 1$ length alternative and uses that same witness in the Stage 13.6 wrapper. Lines 212–237 give the Bohr-model wrapper; `Proofs13UniformEdgeModels.lean`, lines 109–147 at the primary pin, gives the initial-stage wrapper.

The dated span-localization source

`Proofs13SpanLocalization.lean`, lines 66–120, supplies the already existing linear-span parent localization and its large-scale-qualified square theorem. The subsequent fixed-density and uniform-density square sources complete the dated comparison.

The complete fixed-density square construction

`Proofs13CompleteSquareExtraction.lean`, lines 13–48, gives the exponent and the complete context-to-square theorem for every fixed exact density $0 < \alpha \leq 1$. Lines 49–90 assemble its proof; lines 94–104 retain the earlier $\alpha \leq 1/6$ statement as a specialization. Section II.210.5 compares its quantifiers and bound without claiming a new build or formal verification.

The final uniform actual-density square theorem

`Proofs13UniformSquareExtraction.lean`, lines 12–24, states one threshold for all contexts above a fixed positive lower density bound. Lines 25–85 assemble its proof. This final source amendment resolves the older fixed-exact-density quantifier limitation while leaving the arbitrary-retained-witness obstruction unchanged.

The frozen inventory additionally includes the following files, so that the surrounding scalar and initial-stage interfaces can be inspected without a network fetch:

- `Proofs13ExponentMargins.lean`
- `Proofs13IntegerBudgets.lean`
- `Proofs13LargeScaleBudgets.lean`
- `Proofs13UniformEdgeModels.lean`

Their inclusion is provenance, not an additional claimed audit of every theorem they contain.

The comparison from the primary pin to the newer pin adds the common-step-cover file and updates its import and status records. It does not change the stage, grid, or budget definitions used here. In particular, the stages file retains Git blob

`1de556b791c7dac188d91f554231b11e2cf31cea.`

The public source inventory is curated: it contains eighteen relevant Lean snapshots covering the primary full-budget definitions, conditional cover and parent localizations, and the imported coefficient and short-parent APIs, with a repository path, full commit, byte length, SHA-256 digest, Git blob digest, and pinned URL for each entry. Matching recorded hashes identify the unchanged stage, grid, and common-step-cover files at the aligned-parent pin, and the unchanged stage, grid, and budget files at the 77f7 fixed-density complete-square pin. Raw retrieval responses, failed-fetch logs, status inventories, and third-party paper copies are not part of the public package.

The companion checks the bundled raw source bytes against this curated inventory. Git identities are computed using the exact formula $\text{SHA1}(\text{blob} \parallel \text{decimal}(|B|) \parallel \text{NUL} \parallel B)$ for source bytes B , with no text or line-ending normalization. The record fixes the inspected source version; it is not an independent signature of authorship or a Lean proof check. No literature-priority claim, exhaustive repository novelty claim, full-project Lean verification, or new global quantitative Szemerédi bound is made.

V.198 Repairs and questions left open

Source 52, delivered Section 10.

The examples determine several necessary distinctions, but do not supply a complete replacement theorem over composite moduli.

Retain the prime hypothesis. This is already the corrected source statement and the scope of the pinned theorem `lemma_13_7_holds`. No new result is needed to see why our examples do not apply there: the moduli in every construction are nonprime. A fresh formal audit of the corrected theorem is outside this report.

A unit common input step is insufficient. Section II.216 disproves this proposed repair, even with genuine initial-stage predicates if no upper bound on R is added. The dense row domain occupies one residue class inside a unit-step interval. Making the containing progression parameter invertible does not by itself force an ambient affine extension of a Freiman map on its sparse row subset.

The simultaneous geometric strengthening remains a separate question.

Theorem II.217.1 rules out repairing the all-moduli implication merely by adding genuine initial

stages, the exact cutoff, $\ell R \leq \ell Q$, or all six numerical budgets together. Its common step is a nonunit. The combined interface that adds a unit common step to that full package remains a separate question: the unit-step examples in Section II.216 do not supply it. No positive theorem or counterexample for that simultaneous package is claimed here.

Coordinate-affine conclusions change the problem. The map on T is exactly $j \mapsto j$ in the progression coordinate $2j$. A conclusion phrased in such coordinates avoids this particular nonextendability demand. But its compatibility with later stages, ambient coefficient extraction, translations, and bilinear conclusions must then be proved. The current predicates cannot simply be read as if they already used this weaker meaning.

Torsion compatibility offers a concrete test. Lemma II.219.2 supplies a necessary two-point test for every proposed row slope. For an entire row, pairwise congruence solvability alone is not asserted to be sufficient for one common slope; compatibility of all congruences must also be checked. A useful replacement theorem might retain a slope in a module or a compatible homomorphism on a difference subgroup, together with explicit extension information. This is a proposed direction, not a theorem established here.

Formalize the negative statement separately. A Lean development could construct the context and records and prove negations of the two old `Prop` definitions without touching the corrected theorem. The modular unit lemma, Freiman no-wrap argument, arrangement bijection, and real-power strict inequality are natural components. This report supplies written proofs and finite exact diagnostics, not such a formal certificate. In particular passing the companion does not establish the universal real inequalities or compile any Lean file.

[write] **7 October 2026; sources 47–53.** Not done at HEAD 06b30ca52: the repository formalizes no negation of the two preserved Stage 13.7 encodings (contrast `lemma_16_10`, whose negation was formalized after source 50).

V.199 Reproducibility and verification limits

Source 52, delivered Section 11.

The accompanying archive contains this editable TeX source, the PDF, a standard-library Python companion, tests, a deterministic builder, and five curated Lean source snapshots with JSON provenance. It does not contain full third-party papers, raw service receipts, private review notes, or a repository checkout. No repository files were changed or uploaded.

V.199.1 What the exact companion checks

The bounded checks cover:

1. the integer exponents in (II.213.4), the two explicit length thresholds, and the symbolic parameter identities in (II.215.1);
2. the full-budget cutoff, spectrum and radius exponents, strict symbolic margin $\beta < \sigma$, and finite toy floor inequalities;

3. the rational brackets $5 < 56/11$ and $16/3 < 6$, which combine with the stated elementary bounds on π to verify the floor;
4. exact small-model Freiman relations using an eightfold convolution indexed jointly by domain sum and image sum, with repeated entries allowed;
5. the one-point and residue-class agreement obstructions on bounded prime-power examples;
6. exact spectral alias sets, arrangement-energy normalization, and the two-point $\mathbb{Z}/18\mathbb{Z}$ model including its failure at order nine;
7. bounded examples of the congruence-solvability criterion and the symbolic unit-step mass ratios;
8. byte lengths, SHA-256 hashes and Git blob SHA-1 hashes of the five fixed source snapshots.

These finite checks help expose indexing, sign-independent spectral, exponent, and normalization mistakes. They do not numerically inspect the enormous witnesses, prove the analytic Fourier bound by floating point, or replace the written arbitrary-parameter proof. Large constants are represented by literal formulas and small exponents; the companion does not construct U, k, N from the genuine initial-stage example or evaluate the full exponent towers.

V.199.2 Build and packaging contract

From a fresh extracted Report280 directory, run

```
python3 -I -B build.py verify
python3 -I -B build.py reproduce -output /tmp/report280-rebuild
```

The output path must be fresh or empty, absolute, outside the source directory, and outside protected system locations. The Linux builder checks a closed file inventory and rejects symlinks, hard-link aliases and nonregular files. It snapshots source bytes; stages checks in a separate directory; runs all tests and the companion normally and under `-O`; checks identical standard output; uses bounded subprocess logs and timeouts; and verifies that original and staged source inventories remain unchanged.

The PDF is compiled twice with shell escape disabled and deterministic metadata. Undefined references, missing citations, overfull boxes, and unresolved rerun requirements are build failures. Archive members, permissions, ordering, timestamps and compression settings are fixed. An internal SHA-256 manifest covers every public payload file, while an external SHA-256 sidecar identifies the actual ZIP bytes. Neither a hash nor deterministic packaging proves the mathematics or authenticates its author.

A successful replay from the distributed package also requires the rebuilt PDF to match its distributed bytes. This is a toolchain-sensitive requirement, not a promise of byte identity across every TeX distribution. The README records the reference tools and alternatives for interpreting a mismatch. Rendering and visual inspection are separate quality checks performed before release; the builder's text-level TeX gates are not a replacement for them.

V.200 Pinned source map

Source 52, delivered Appendix A.

All five source files below are included under `provenance/sources`. Full digest values and verified URLs are in `provenance/source_manifest.json` and `SOURCES.md`. The common source commit is the 40-character pin stated in Section F.57.

1. [Definitions.lean](#). Progressions and properness: lines 82–94. Unnormalized Fourier transform: lines 135–137. Ambient affine agreement: lines 227–229. Centered absolute value: lines 277–279. Standard Freiman predicate: lines 311–314. Additive tuples: lines 316–330.
2. [Sections12_13.lean](#). Arrangements and counts: lines 60–99. Edge functions, sections and separate Freiman condition: lines 198–247. Context: lines 304–314. Initial-stage constants, records and predicates: lines 333–414. Stage 13.6: lines 415–464. Stage 13.7 records, predicates and both targets: lines 471–515. Prime-corrected proposition: lines 517–525.
3. [Section10.lean](#). Exact spectral cutoff, spectrum ceiling and corrected Bohr radius: lines 434–454.
4. [Proofs13BilinearExtraction.lean](#). Prime-qualified extraction theorem and six numerical assumptions: lines 16–32. The report checks these numerical conditions without asserting the prime hypothesis.
5. [Proofs13CompleteRowExtraction.lean](#). Prime-modulus affine base partition: lines 14–42. Corrected theorem `lemma_13_7_holds`: lines 46–53. This source is included to keep the distinction between the refuted definitions and the corrected theorem directly reviewable.

The snapshots are a curated reading set, not a standalone Lean project. Their imports are not bundled, and the archive does not claim an offline Lean compilation. The mathematical proof depends on the displayed definitions; the corrected theorem file documents the scope boundary.

V.201 A concise reconstruction guide

Source 52, delivered Appendix B.

For a manual check of the counterexample, the following order avoids both giant-integer computation and dependence on the optional finite tests.

1. Read `LinearOn` and the two target quantifiers, then the complete Stage 13.6 and weak Stage 13.7 predicates.
2. Verify the strip context by the strict bound $8(m - 1) < N/2$ and the exact identity $C(h) = N^{16}E_8(T)$.
3. Check the two Stage 13.6 masses against the number of critical heights. For step two their ratios are δ, λ ; for the unit-step enlargement they are $2\delta, 2\lambda$.
4. Prove the odd multiplier $2a - 1$ is a unit. This forces at most one agreement on each parity class.

5. Cancel the positive number of rows in the first output mass inequality. Compare the resulting upper bound 2^{939} or 2^{940} with the exact output length lower bound.
6. For genuine initial data, check both spectral aliases, the exact floor $\lfloor 16/\pi \rfloor = 5$, and the Stage 13.5 coefficient estimate appropriate to the chosen step.
7. For the full numerical package, check $\beta < \sigma < uw$, all nine scale conditions, and the separate lower-target width budget. Retain the distinction between its nonunit step and the unit-step examples with fewer simultaneous extra conditions.

This order isolates the obstruction from the large constants: the algebra is elementary, while the giant scale is used only to make a positive-power length lower bound exceed a fixed row-density cap.

V.202 Scope and further questions

Source 53, delivered Section 9.

The result is a quantitative consequence of the source constructions on the full density interval. It preserves actual-density quantities uniformly and exhibits an explicit closed threshold. The seventh-term dominance is exact for the displayed seven-term budget; no claim is made that this is the smallest possible threshold for square extraction.

The basic uniform source theorem remains fully credited. The geometric route here strengthens selected outputs by preserving $\ell R \in \{m, m+1\}$ and using a new localized height progression. It does not repair a proposition that asks for a grid from every arbitrary witness satisfying a weaker retained interface, and does not compare the new square with the length of an independently chosen old Stage 13.9 progression.

The following questions delimit useful further work:

1. Can the existing Stage 13.6 wrappers formally export $\ell R \leq m+1$, and can the short-carrier and balanced-parent arguments be implemented with unchanged carriers under reversal? The proof-level construction above identifies precisely the extra information to retain.
2. Can the threshold be reduced by deriving only the width and short-parent budgets needed for the sharper route? We retained the stronger sixth budget that also pays all scales of the existing square API; its economy is not asserted.
3. For a chosen constant $\lambda < 4$, what is the best explicit threshold for $k \geq N^{\lambda\gamma(a)}$ uniformly? The coefficient-bearing result (II.202.1) is a suitable starting point, but this report only proves the stated clean value three at (F.58.5).
4. How should this actual-density refinement be composed with subsequent Fourier and global iteration interfaces while preserving their hypotheses and printed target exponents? No such global conclusion is derived here.

These are questions about further quantitative or formal developments, not assertions that every current upstream extraction route requires them.

V.203 What the reproducible companion verifies

Source 53, delivered Section 10.

The proof is symbolic at the full scale. No natural number near $N_0(a_0)$ is constructed, no enormous prime is searched for, and no denominator of astronomical bit length is allocated.

The standard-library Python companion checks bounded exact exponent identities and coefficient inequalities underlying the scalar estimates, the natural floor/half and ceiling budgets, short rectification on finite models, balanced aligned-parent partitions, common-step cover capacity and side-fit arithmetic. It uses integers and rational numbers, not floating-point surrogates for the thresholds. These finite checks detect transcription and indexing mistakes; they do not prove the universally quantified analytic inputs or replace the proofs in the article.

The companion also verifies the curated source files against their byte counts, SHA-256 digests and Git blob identities. The latter use

$$\text{SHA1}(\text{blob} \parallel \text{decimal byte length} \parallel \text{NUL} \parallel \text{file bytes}).$$

Matching hashes identify the inspected source bytes, not their logical correctness or the health of a whole dependency build. No Lean source was edited, no new Lean proof was produced, and no kernel replay or axiom-dependency audit is claimed.

The ZIP contains the editable L^AT_EX source, PDF, bounded companion, tests, offline builder, and twenty-three relevant raw Lean source snapshots with a curated JSON inventory. Private reviews, raw connector receipts, unrelated source inventories and full third-party papers are excluded. The builder runs diagnostics and tests normally and with Python optimization, rejects source mutation, uses a fixed PDF build epoch and deterministic archive metadata, and supports replay from an extracted delivered archive. The README states its platform, toolchain and output restrictions. A manifest detects byte changes; it is not an independent trust anchor for an archive controlled by an adversary.

V.204 Pinned source map

Source 53, delivered Appendix A.

All paths below are relative to `Combinatorics/Ramsey/Lean/GowersSzemerédi/` in [VladimirReshetnikov/ProveIt](#), at the frozen comparison commit

$$\text{bf2fe146d323cf2828ca0de63c9a965ca0638eaa}.$$

This comparison was checked at 01:17 UTC on 7 October 2026. Previously inspected files were unchanged through that pin; the common-step and coefficient files were additionally fetched and Git-blob-checked at the pin itself. Each public record states both the effective comparison commit and the originally observed commit. This is a dated source comparison, not a claim about an unpinned moving branch.

Context and scalar definitions

[Sections12_13.lean](#), lines 303–314, gives the exact context; the surrounding file gives the arrangement counts and successive stage predicates. [Section10.lean](#), lines 434–455, defines the Section 10 cutoff and corrected rounded radius. [Section05.lean](#), lines 44–63, defines the recurrence threshold whose quadratic specialization is (II.197.9).

Initial progression and scaled recurrence

[Proofs13InitialProgression.lean](#) supplies `lemma_13_4_holds`. [Proofs13QuadraticPartition.lean](#), lines 76–84, gives the scaled Stage 13.5 constructor. [Proofs13LargeScaleRecurrence.lean](#), lines 33–36, defines the three recurrence scales; lines 112–120 verify cutoff admissibility.

Monotonicity and the uniform parameter margin

[Proofs13UniformDensityParameters.lean](#) proves the relevant monotonicities and the factor-two recurrence margin. [Proofs13ExponentMargins.lean](#) supplies the spectrum-bound comparison used in that margin. [Proofs13UniformEdgeModels.lean](#), lines 22–50, treats monotonicity of the rounded Section 10 radius including its denominator.

All-density models and the short witness

[Proofs13DenseEdgeModels.lean](#), lines 93–174, supplies all-density height models. [Proofs13AllDensityExtraction.lean](#), lines 13–49, uses them in the initial-stage wrapper. [Proofs13CommonStepSelection.lean](#), lines 92–181, constructs the actual m or $m + 1$ witness and the Stage 13.6 record. Its lines 212–236 pass through the same construction for the Bohr-model wrapper.

Uniform recurrence and rounded budgets already upstream

[Proofs13UniformDensityRecurrence.lean](#), [Proofs13UniformRoundedBudgets.lean](#), and [Proofs13UniformDensityBudgets.lean](#) prove uniform Stage 13.5 and 13.6 existence with actual-density parameters. The explicit estimates here pay those underlying budgets directly.

Rows and all-scale coefficient localization

[Proofs13CompleteRowExtraction.lean](#) proves the prime-corrected Stage 13.7 theorem. [Proofs13CoefficientExtraction.lean](#), lines 20–96, constructs the width-qualified Stage 13.8 output. [Proofs13CoefficientPartition.lean](#), lines 84–137, contains the all-scale localization and the exact exponent identity used in (II.201.2).

Existing square extraction and its current uniform theorem

[Proofs13SquareExtraction.lean](#), lines 74–130, provides indexed-span square extraction with exact scale assumptions, actual-density mass and the floor/half side bound. [Proofs13SquareScale.lean](#), lines 11–22, verifies the natural floor/half estimate. [Proofs13CompleteSquareExtraction.lean](#) contains the fixed-exact-density complete theorem and the exponent definition. [Proofs13UniformSquareScale.lean](#) provides monotone transport of the square scales. [Proofs13UniformSquareExtraction.lean](#), lines 15–24, already states the density-uniform theorem; its proof obtains actual-density mass before weakening it to the fixed lower density.

The last uniform theorem was already present at commit

4f8cd11982e6037060b594a3da4296260d26f715.

It is unchanged at the frozen comparison pin. The intervening update adds explicit Fourier-exponent, odd-side trimming and Fourier-box-transport files, but does not supply the explicit threshold or the sharper selected-witness result proved here. No novelty-priority conclusion is inferred from this finite source comparison.

V.205 Repository integration and proof dependencies

Source 54, delivered Section 10.

V.205.1 A precise replacement, not an untracked global claim

The interface is pinned to repository commit 6a6d7961c3c261da7dc18def13fff30b12dd5c04 [223]. The principal source paths, relative to `Combinatorics/Ramsey/`, are

```
Papers/sz-thm-gowers-proof/sz-thm-gowers-proof.tex
Lean/GowersSzemerédi/Sections08_09.lean
Lean/GowersSzemerédi/Proofs08QuadraticBias.lean
Research/GowersSzemerédi/local-quantitative-refinements/README.md
```

[write] **7 October 2026; sources 54–59.** *Checked at HEAD fe7165a3c.* The four paths exist; `Sections08_09.lean` has the blob read at the pin, `proposition_8_1` is at its line 29 and `proposition_8_1_holds` at line 1069 of `Proofs08QuadraticBias.lean`. The theorem was not added beside the formal proposition: by its placement (7dc1ee29d) this source is source 54 of this report, its labels carry the prefix `gsr:fm:54:`, `gsr:en:54:` or `gsr:fz:54:`, and its files are shipped under the prefix `54-single-phase-`; the formalization plan named below is shipped as `54-single-phase-FORMALIZATION.md`.

The edited transcription and the statement module were read at the pinned commit. The proof module was located through the connected repository search; its statement and odd-modulus interface were checked, not its entire implementation. The existing source audit for derivative-spectrum symmetry was inspected to distinguish that global multilinear topic from the present localized Section 8 result.

The new theorem should first be added beside the existing proposition, not substituted silently into an aggregate research document. A useful formal layering is: finite Fourier product identity; weighted common-frequency extraction; cyclic quadratic untwisting; balanced amplitude specialization; graph certificate and cycle stability. The first four layers imply the old proposition by taking $\psi_s = \psi$ and weakening the numerical bound. The last layer is independent of the density-increment iteration.

V.205.2 Dependency ledger

The dependency structure is small enough to state without a diagram. Lemma II.268.1 depends only on finite Fourier inversion. Theorem II.268.2 adds the triangle inequality and Parseval. Theorem F.59.1 adds one quadratic identity and the exact amplitude of a balanced indicator. The all-group extension adds Proposition II.270.2.

For the extremal theory, Lemma II.271.1 and the elementary certificate of Lemma II.271.2 give all single-offset ceilings and equality cases. Theorem II.272.1 retains the deficit in that certificate; Lemma II.272.2 converts it to function distance. The simultaneous-intercept result is the one place where a named classical graph theorem is invoked, and its extremal-value proof is included.

No large inverse theorem, Freiman theorem, asymptotic input, probabilistic existence theorem, or unproved repository statement is needed for these proofs. Conversely, the present results do not construct the affine selector supplied to Section 8. The quantitative cost of obtaining that selector, and the costs of subsequent phase flattening and iteration, remain part of any global application.

V.205.3 Status of verification

The accompanying checks are reproducible but are not a substitute for the proofs. The exact graph tests enumerate **31,030** rational probability distributions, using integer arithmetic for the triangle-free certificate and testing **227** equality instances against the cycle classification. The exact four-frequency obstruction is checked at five rational values of t .

Independent floating-point calculations check **2,016** mixed weighted factorization cases, **2,540** offset-power and ceiling cases, **402** weighted balanced-function cases, and **1,440** constructive stability cases. The largest recorded factorization or global power-identity discrepancy is below $1.9 \cdot 10^{-15}$, with an acceptance tolerance of 10^{-10} . The backward/forward chirp-sign identity is checked separately, and even cyclic refinements and Example II.270.3 are tested. In addition, **1,440** short-orbit repairs and **7,704** exact rational polarization identities on product groups are checked, together with coordinate periodicity. The random seed is 20261006. Full counts and environment versions are recorded in `data/verification.json`.

No Lean proof is bundled or claimed checked. The companion formalization plan specifies candidate statements and proof boundaries, rather than presenting uncompiled Lean syntax as a verified artifact. The mathematical proofs and source comparisons have been self-audited, not independently peer-reviewed.

V.206 Further research questions

Source 54, delivered Section 11.

The following problems distinguish missing improvements in the present theorems from harder changes to the global inverse argument. None is being asserted to be a previously published open problem.

Question V.206.1 (The best balanced-indicator coefficient). *For fixed odd N and density δ , determine the largest coefficient replacing $1/(B\sigma)$ in Theorem F.59.1, uniformly in P, λ, μ after normalizing by the actual input energy. Does the balanced constraint yield a strict improvement over the amplitude-only bound, uniformly as $N \rightarrow \infty$? A solution must respect that the elementary coefficient-one extremizer need not be a balanced indicator.*

Question V.206.2 (The exact stability modulus). *For cycles of order $m \geq 5$, determine the optimal function $\Phi_m(\varepsilon)$ bounding distance to the extremal star family. Is there a common asymptotic leading coefficient for all $m \geq 5$? The four-frequency family gives a lower bound, while (II.272.1) supplies a dimension-free upper bound of the correct exponent but not a matching coefficient.*

Question V.206.3 (Pointwise-constrained repair). *Suppose $|g| \leq 1$ and the offset energy is almost extremal at a prescribed value of $\|g\|_2$. Can one repair g to an extremizer while preserving the pointwise cap, or quantify the unavoidable increase in the cap? The phase-preserving Fourier repair in Lemma II.272.2 does not answer this. Balanced indicators are an even more restrictive version.*

Question V.206.4 (Sharp localized ceilings). *For a prescribed interval or Bohr set P , compute the maximum of $\mathbb{E}_{h \in P} |A_\eta(g)(h)|^2 / \|g\|_2^4$ under specified amplitude constraints. The coverage estimate (II.274.2) is deliberately coarse. The phases of Fourier coefficients, absent from the full-group identity, now matter through the Fourier transform of $\mathbf{1}_P$.*

Question V.206.5 (Uniform extraction from a family of windows). *Find useful conditions on a family of nonnegative windows w_j under which one common frequency attains a fixed fraction of each window's individual optimal score. Theorem II.268.2 selects a common phase across translates of one window, not simultaneously across unrelated windows. Aggregate extraction is easy; per-window control is the substantive problem.*

Question V.206.6 (Higher-degree common-phase extraction). *When a higher-order derivative selector has an exact polynomial primitive, which local hypotheses force one global polynomial phase to replace the translate-dependent phases in a higher-order analogue of Proposition 8.1? The obstruction to simply repeating the present argument is that the relevant lower-order Fourier coefficient becomes a higher-order correlation, rather than an ordinary Fourier product with controlled ℓ^1 norm.*

Question V.206.7 (Non-symmetric selectors on a restricted shift set). *For a nonsymmetric bilinear selector, can local extraction be organized using its alternating radical and projective Fourier blocks while retaining a useful common structured object? Global symmetry thresholds are already treated in other repository research; the distinct issue is a weighted shift set and a local window conclusion without silently assuming integrability.*

Question V.206.8 (Several intercepts and graph stability). *Which Cayley graphs Γ_S admit dimension-free power-spectrum stability near the clique ceiling, and with what exponents? The transition between four-cycles and longer cycles shows that the extremizer space, not just the clique number, is relevant. A classification should avoid assuming that every maximizer is supported on a single clique.*

Question V.206.9 (A faster certified selector). *Can the $O(N^2 \log N)$ all-frequency scan be reduced when the overlap spectrum is sparse, or when the window is an interval, while keeping an a posteriori certificate for the retained score? Randomized sampling from the overlap spectrum is natural, but a useful running-time theorem needs an explicit lower bound for success probability or a suitable stopping certificate.*

Question V.206.10 (Quantitative propagation through Section 8). *Combine the common-phase theorem with the strongest compatible selector construction and phase-flat partition in the repository. Track every normalization and exceptional-set loss to determine which downstream constants improve. An improved local coefficient alone does not establish a better exponent in a final bound for $r_4(N)$ or $r_k(N)$.*

Question V.206.11 (Formal generality of quadratic refinements). *Develop a reusable finite-abelian-group quadratic-refinement API, with classical cyclic refinements as a special case and doubled roots for even moduli. The conceptual theorem is elementary, but formalizing the coordinate-independence calculation once would avoid repeated incorrect divisions by two in local inverse steps.*

V.207 Conclusion

Source 54, delivered Section 12.

Proposition 8.1 admits a stronger conclusion than its original translate-by-translate extraction: one quadratic phase, with the prescribed quadratic coefficient, suffices across the entire translation family. The exact factorization retains a spectral overlap and an amplitude weight; both sharpen

the lower bound. Quadratic refinements identify the correct all-group extension and its even-order limitation.

The same Fourier representation resolves the normalized intercept energy problem completely. Its equality cases change at small character orders, and its optimal recovery exponent changes between orders four and five. The four-frequency obstruction explains why a superficially plausible square-root repair theorem fails in the latter range. These statements give explicit local improvements and testable structural interfaces for the ProveIt Ramsey development without overstating their global consequences.

V.208 Further research questions

Source 55, delivered Section 12.

The following are proposed problems arising from the proved results. They are not all claims of previously documented open conjectures. The first four are particularly direct continuations of the present work.

Question V.208.1 (The cubic coefficient). *Determine the optimal coefficient of e^3 in $\mathfrak{S}_{d,G}(e)$. Is it determined by finitely many torsion statistics of $\widehat{G}$ for fixed d ? Does it depend on the smallest available odd character order even when the quadratic coefficient does not?*

The proof explains why only order-two torsion can affect the quadratic coefficient. At the next order, six-vertex configurations and the next-order deformation of the extremizing character modes both enter. One must optimize those deformations, not merely expand the families (II.243.2) and (II.243.6).

Question V.208.2 (The exact even-order curve for $d \geq 3$). *For some group-independent interval $0 \leq e \leq e'_d$, is $\mathfrak{S}_{d,G}(e)$ equal to (II.243.5) on every even-order finite abelian group? If not, determine the first coefficient at which a multi-mode deformation improves the one-character family.*

The equality is proved for $d = 2$ in proposition II.245.1. The present higher-degree argument proves only its second-order optimality.

Question V.208.3 (Quantitative rigidity). *Give explicit rates in theorem II.237.3 in terms of the second-order deficit $\frac{2}{q}\varepsilon + b_d(G)\varepsilon^2 - D_d(f)$ and ε^3 . Can one obtain a stability estimate for the full phase perturbation, rather than just its normalized tangent?*

The proof already isolates three nonnegative losses: amplitude deficit, two-torsion profile deficit, and failure of Fourier mass concentration. Keeping quantitative versions of all three is a plausible route.

Question V.208.4 (An explicit entry threshold). *Extract a usable numerical e_d from the finite-group proof of the Eisner–Tao entry theorem, then combine it with theorem II.242.1. Can the resulting threshold and the local radius be made singly exponential in d , rather than involving a coarse expansion in 2^d vertices?*

This is an optimization of an effective published argument, not a missing logical hypothesis of theorem II.237.1.

Question V.208.5 (Intervals and boundary terms). *Find the sharp analogue of the distance expansion for the normalized interval norm $U^d([N])$, uniformly in sufficiently large N . Which coefficients reflect the boundary-conditioned cube distribution, and what additional dependence on N is unavoidable?*

The interval near-extremizer theorem in [7] supplies an entry point, but the three-wise independence used here is absent. A direct transfer of the finite-group constants would therefore be unjustified.

Question V.208.6 (Mixed cubical near-extremizers). *For a tuple (f_ω) of bounded functions whose mixed cube average is near one, determine the sharp joint distance to a compatible tuple of phase polynomials. Is there a finite-dimensional spectral problem replacing (II.241.7)?*

Here compatibility matters: individually good phase approximations need not satisfy the exact mixed cube identity. An analysis of simultaneous phase gauges would be needed before a scalar expansion can be used.

Question V.208.7 (Approximately independent cube samplings). *Extend the local theorem to sampling laws for which every triple of vertex evaluations is close to independent, with an explicit error in terms of the triple marginal defects and the four-point circuit law.*

Such a result could separate the analytic cancellation from the exact group geometry. It would also clarify which weighted or local cube models admit the same leading coefficient.

Question V.208.8 (Compact groups and spectral type). *Develop the same sharp theory for second-countable compact abelian groups, replacing parity by the existence of a nontrivial order-two character. Can the tangent theorem be formulated uniformly without finiteness of the phase space modulo constants?*

Most local calculations extend through Haar measure and Fourier series; existence or replacement of a minimizing phase requires an explicit argument. No compact-group theorem beyond the imported entry statement is asserted here.

Question V.208.9 (Noncommutative endpoint constants). *For the Pauli uniformity norms of unitary matrices, determine the sharp leading and second-order distance constants near an exact extremizer. Which part of the scalar two-torsion dichotomy survives, and which is replaced by commutator data?*

Bao, Briët, Castro-Silva, van Dordrecht, and Helsen prove a fourth-order near-extremizer theorem in this noncommutative setting [226]. Their result provides a concrete comparison problem. The commutative Fourier circuit calculation here does not transfer unchanged, and no noncommutative theorem is claimed in this article.

Question V.208.10 (Certified integration). *Formalize the four-vertex classification, the six-triple Hölder estimate, and the explicit local theorem in Lean. Can the phase-entry input be kept as a clearly named interface while the local result is verified independently of the full Szemerédi development?*

This is a modular target: the local theorem requires finite-group averages, character orthogonality, Hölder, and polynomial inequalities, not the entire inverse theorem.

V.209 Validation, provenance, and integration

Source 55, delivered Section 13.

V.209.1 What was checked computationally

The supplied `code/verify.py` has no network dependency and uses a fixed seed. Its executed output is in `data/verification_results.json`. It checks the following items.

Check	Number
Four-vertex subsets, dimensions 2 through 5	37,851
Exact minor-gcd classifications, dimensions 2 through 4	1,891
Exact rational cube and moment identities	63
Bounded complex test functions for the local inequalities	420
Of these, functions satisfying the explicit local radius	159

The exact checks include real quartic coefficients and the second and fourth moments of signed cube differences. The complex checks include the Fourier profile, amplitude bound, aggregate cubic remainder, master inequality, and, when the radius hypothesis holds, its explicit inversion. They run on cyclic groups and selected products, including odd, even, and noncyclic examples. High-precision checks on the even family and the odd family over $\mathbb{Z}/3\mathbb{Z}$ confirm the limiting quadratic coefficients for $d = 2, 3, 4, 5$.

For example, at $d = 3$ and $t = 10^{-4}$, the ratio

$$\frac{D_d(F_t) - \varepsilon(F_t)/4}{\varepsilon(F_t)^2}$$

is approximately 0.2031250108593753 for the odd family, compared with $13/64 = 0.203125$, and 0.3281250278906264 for the even family, compared with $21/64 = 0.328125$. These numbers are illustrations of the exact Taylor calculations, not evidence substituted for proof.

The full-interval U^2 calibration is checked separately: 1,024 bounded complex functions are generated, of which 859 have defect at most $1/2$, and 255 points on the even-order attaining curves are checked, including the endpoint. The absolute tolerance there is $2 \cdot 10^{-7}$ because the inverse formula has a square-root singularity at defect $1/2$.

Floating checks use a stated absolute tolerance; rational checks use integer or `Fraction` arithmetic. No random test verifies the qualitative entry theorem, the all-group quantifiers, or the asymptotic rigidity theorem. Those are justified by the cited theorem and the mathematical proofs above.

V.209.2 Suggested repository placement

A self-contained directory beneath

`Combinatorics/Ramsey/Research/GowersSzemerédi/`

with a name such as `polynomial-phase-endpoint-stability` avoids collisions with the numbered source files in the existing merged report. The article is independent of that report’s macro system and bibliography. It can also be integrated thematically, with labels prefixed to prevent collisions. The included `FORMALIZATION_PLAN.md` describes the logical layers and candidate theorem interfaces; it contains no purportedly compiled Lean code.

[write] **7 October 2026; sources 54–59.** The suggested directory was not used: by its placement (7dc1ee29d) this source is source 55 of this report, its labels carry the prefix `gsr:fm:55:`, `gsr:en:55:` or

gsr:fz:55:, and its files are shipped under the prefix 55-endpoint-stability- (the formalization plan as 55-endpoint-stability-FORMALIZATION_PLAN.md, the verifier as code/55-endpoint-stability-verify.py).

The inspected source reference is commit 128f514afce0bf7eb48133fb30dd3f8e034ea3a2. The repository was moving during inspection; this identifier records a reproducible comparison point, not an assertion that no later material exists. In particular, a declaration of type **Prop** in the Gowers catalogue must not be mistaken for a proof of that declaration.

V.209.3 Logical dependency summary

The four-vertex classification, Fourier inequalities, cancellation, explicit local theorem, matching examples, phase-class separation, and rigidity arguments are proved in this article. The sole substantial external entry result for the global near-endpoint theorem is theorem II.243.1. Standard finite Fourier analysis and Hölder are used in fully displayed forms. The exact U^2 formula is included as an elementary calibration rather than a priority claim.

No improved global Szemerédi bound, no resolution of the repository’s open general inverse/integration interfaces, and no proof-assistant verification is claimed. The mathematical output is a sharp quantitative endpoint theorem and its local and structural refinements.

V.210 A compact coefficient ledger

Source 55, delivered Appendix A.

For convenient independent checking, the constants used in the proof are collected here:

$$\begin{aligned} q &= 2^d, & C_d &= (6^d - 2q^2 + q)/8, \\ A_d &= q(q-1)(q-2)/24, & J_d &= A_d - C_d, \\ B_d &= 10\binom{q}{5} + \sum_{s=6}^q \binom{q}{s} 2^{s-6}, & K_{0,d} &= q(q-1)/8. \end{aligned}$$

The two exact simplifications driving the main theorem are

$$\frac{8}{q^3} \left(K_{0,d} + \frac{C_d}{2} \right) = \frac{3^d - 1}{2q^2}, \quad \frac{8}{q^3} (K_{0,d} + A_d) = \frac{q^2 - 1}{3q^2}.$$

For $d = 2, 3, 4$ the coarse local constants are

d	B_d	K_*	H_*	R_d
2	0	5/2	1	121/32
3	608	21	620	25813/256
4	713568	170	713668	51249029/4096

The quadratic coefficients are optimal; these cubic upper constants are not. The five-factor cancellation, rather than the numerical size of B_d , is what establishes the $O_d(e^3)$ scale.

V.211 Why a purely formal Taylor expansion is insufficient

Source 55, delivered Appendix B.

There are three distinct uniformity issues in this problem. First, $D \rightarrow 0$ does not imply $\|F - 1\|_\infty \rightarrow 0$ when the group is allowed to vary. A perturbation may remain large on a set of vanishing relative size. Consequently one cannot simply write $F = e^{ib}$ with $\|b\|_\infty \rightarrow 0$ for an arbitrary near-extremizer and expand the exponential. The upper theorem instead expands the finite cube product exactly and bounds its tail using triple independence and the disk constraint. Second, the distance in (F.60.2) is an infimum over all phase polynomials, not just over constants. Computing a family's mean gives a distance upper bound unless other phases are excluded. The uniform gap in lemma II.243.2 supplies precisely that missing step for the sharpness examples. Third, a local inequality around an already known phase does not imply a global near-extremizer theorem without an entry argument. The explicit separation of theorem II.242.1 from the Eisner–Tao input makes this dependency visible. None of the finite tests in section V.209 replaces it. These distinctions are also useful for formalization: the local analytic estimate, the discrete phase geometry, and the inverse-theorem entry should be separate interfaces rather than hidden assumptions in one statement.

V.212 Deterministic reconstruction and exact diagnostics

Source 56, delivered Section 10.

V.212.1 The decoder

The proof gives the following deterministic reconstruction procedure on an explicitly enumerated domain. Ties may be resolved by any fixed rule.

1. For every $h \in G$, form the weighted histogram of $\phi(x + h) - \phi(x)$ with weights $w(x)w(x + h)$. Let $q(h)$ be a mode, and use zero if the histogram has zero total mass.
2. For every $x \in G$, take the unweighted mode of $q(x + y) - q(y)$ over $y \in G$. Denote it by $L(x)$.
3. Take the weighted mode c of $\phi(x) - L(x)$ on the original domain, with weights $w(x)$. Output $\Phi(x) = L(x) + c$.

Under the theorem's hypotheses the second-stage modes are unique and define a homomorphism. The first-stage ties need not be unique, but the final repaired affine map is unique, so every tie-breaking rule gives the same answer. Off-support values of ϕ do not affect the output.

With hashable group elements and exact addition, equality, and weight arithmetic, the algorithm uses $O(N^2)$ group and dictionary operations. This bound does not count bit complexity of rational arithmetic as constant; large denominators incur their usual cost. Histograms can be processed one shift or one point at a time, so the algorithm itself does not require storing all N^2 pairs. The supplied implementation enumerates both groups as products of cyclic groups and uses exact rational weights.

The hypothesis can also be checked without numerical Fourier transforms: compute $r(h)$ and then use

$$u^4 = \mathbb{E}_h r(h)^2 - \alpha^4.$$

The same first-stage histograms give $\mathcal{Q}(w)$ and $\mathcal{D}_w(\phi)$ exactly, by (II.255.2). The returned map can be checked for additivity on all N^2 pairs independently of the proof.

V.212.2 Verification scope and results

The file `code/repair.py` is a reusable reference decoder. The file `code/verify.py` performs integer and rational checks; it uses no floating-point inequality tests. The supplied run completed **33,889 assertions**, including the following families:

- All 1,086 nonempty labeled subsets of $\mathbb{Z}/n\mathbb{Z}$ for $1 \leq n \leq 6$ with two possible labels at each included point. This is an enumeration of three states per point: absent, present with label zero, and present with label one.
- Ninety seeded rational-weight cases on cyclic and noncyclic groups, with weights in $\{0, 1/4, 1/2, 3/4, 1\}$ and labels in $\mathbb{Z}/3\mathbb{Z}$.
- Three nonzero-error weighted examples satisfying the main hypotheses, plus exact closed-form data for a nonzero-error subset example of size 8,192 in a composite-order cyclic group.

The seed is 20261006. The checks include the two energy identities, the exact triangle second moment, the weighted interpolation inequality, the triangle transfer, every pointwise link bound, all six marked-pair bounds, inclusion–exclusion, the self-correction conclusion whenever its hypothesis holds, and the numerical constants in the final bootstrap. Most small arbitrary subsets do *not* satisfy the main uniformity hypothesis; they test the unconditional intermediate statements instead.

An inequality of the form $X \leq A + \sqrt{B}$ is checked exactly by testing $X \leq A$ or $(X - A)^2 \leq B$. This avoids rounding the radicals in the proof-critical estimates. The global proof for every finite group is the mathematical argument in the preceding sections, not an extrapolation from these finite checks.

V.212.3 Two reproducible nonzero-error examples

On $G = \mathbb{Z}/31\mathbb{Z}$, put $w(0) = 1/10000$ and $w(x) = 1$ otherwise, with $\phi(0) = 1 \in \mathbb{Z}/3\mathbb{Z}$ and all other labels zero. The decoder returns the zero affine map and computes exactly

$$\alpha = \frac{300001}{310000}, \quad t = \frac{1}{300001}, \quad \varepsilon = \frac{3480006000000000}{261303480018000000001}.$$

The small weight makes the nonzero label error small without requiring a huge explicitly enumerated domain. The checked uniformity fourth power is

$$u^4 = \frac{29988001799880003}{92352100000000000000},$$

which satisfies $u^4 \leq \alpha^5/1024$.

For an unweighted example let $N = 8193$, $G = \mathbb{Z}/N\mathbb{Z}$, and $A = G \setminus \{0\}$. Give point 1 label one in $\mathbb{Z}/2\mathbb{Z}$ and all other points of A label zero. Then

$$\begin{aligned} \alpha &= \frac{8192}{8193}, & u^4 &= \frac{8192}{8193^4}, \\ t &= \frac{1}{8192}, & \varepsilon &= \frac{5589675}{11451848192}. \end{aligned}$$

These numbers satisfy the hypotheses and the claimed error bound. Here the large example is checked using exact closed-form counts, not by running the quadratic-time decoder on all 8,193 points.

To derive the counts, write $m = N - 1$. The unnormalized pair counts are $R(0) = m$ and $R(h) = N - 2$ for $h \neq 0$. The number of pairs with nonzero label difference is zero at $h = 0$, one at $h = 1, -1$, and two at every other shift. Thus

$$N^3 \mathcal{Q}(\mathbf{1}_A) = (N - 1)^2 + (N - 1)(N - 2)^2,$$

while

$$N^3 \mathcal{D}_A(\phi) = 4(N - 3) + 4(N - 3)(N - 4) = 4(N - 3)^2.$$

The code cross-checks these formulas by full enumeration for $N = 5, 7, 9, 11$ before evaluating the larger instance.

V.212.4 Reproducibility and limits

From the package directory, run

```
python3 code/verify.py --output data/verification.json
python3 code/repair.py
make pdf
```

The mathematical claims do not depend on a numerical solver, a randomized search finding an extremizer, or an unverified external computation. No Lean code is included or claimed to compile. The JSON verification record explicitly labels its checks as finite diagnostics.

V.213 The Gowers interface and a formalization route

Source 56, delivered Section 11.

V.213.1 Which Gowers step is strengthened?

The inspected repository transcription treats Proposition 6.1 as a passage from large Fourier coefficients of many differences to many respected additive quadruples. Lemma 6.2 expresses that count as the additive energy of a function graph. Corollary 7.6 then extracts a restriction that is a Freiman homomorphism of order eight [230, 1]. Our theorem supplies a different, stronger conclusion in a narrower but useful input regime: the domain is Fourier-uniform and nearly all *available* quadruples are respected.

For a map $\phi : A \rightarrow H$, define its finite graph

$$\Gamma_\phi = \{(a, \phi(a)) : a \in A\} \subseteq G \times H.$$

Let $E(A)$ and $E(\Gamma_\phi)$ be unnormalized ordered additive energies. There is an exact identity

$$E(\Gamma_\phi) = (1 - \varepsilon)E(A), \quad E(A) = N^3(\alpha^4 + u^4). \quad (\text{V.213.1})$$

Indeed, graph quadruples correspond bijectively to domain quadruples whose labels satisfy the same relation. No average over H is involved, so this identity also makes sense when H is infinite.

Corollary V.213.1 (A high-agreement alternative to small-restriction extraction). *Assume the subset hypotheses of Corollary F.61.2 and $E(\Gamma_\phi) \geq (1 - \varepsilon)E(A)$ with $0 \leq \varepsilon \leq 2^{-10}$. There is a single subset $B \subseteq A$ of size*

$$|B| \geq (1 - 13\varepsilon/50)|A|$$

on which ϕ is a Freiman homomorphism of every positive integer order, and $\phi|_B$ is the restriction of one affine map on G .

Proof. Let B be the agreement set with the repaired affine map. If $\sum_{i=1}^k x_i = \sum_{i=1}^k y_i$ with all $x_i, y_i \in B$, then applying its linear part and adding k copies of its constant gives $\sum_i \phi(x_i) = \sum_i \phi(y_i)$. The same B works for every k . $\square$

This is not an unconditional improvement to the size guarantee in Corollary 7.6. The normalizations explain why. If

$$\gamma = \frac{E(\Gamma_\phi)}{|A|^3},$$

then (V.213.1) reads

$$\gamma = (1 - \varepsilon)\alpha(1 + \kappa). \quad (\text{V.213.2})$$

For a sparse uniform domain, γ is near α , not near one. Near-maximal energy relative to the domain's *available* quadruples is a different statement from near-maximal energy relative to $|A|^3$. The earlier full-domain/coset energy stability arguments cannot simply be substituted for the relative theorem by ignoring this distinction.

V.213.2 Interaction with the local Section 10 argument

The repository's Lemma 10.9 proves that a locally constructed difference map is a Freiman homomorphism on a specified smaller set; its formal statement includes the nonemptiness needed to avoid vacuous nested assertions [230]. The present proof uses a global triangle identity and an ambient Fourier-uniformity assumption instead. It does not prove a stronger conclusion from precisely the hypotheses of Lemma 10.9. A future integration should expose the relative repair theorem as an additional branch, not overwrite the existing theorem statement.

On a Bohr neighborhood or an arbitrary progression, the needed convolution and sampling identities have boundary and translation-measure issues. Those issues are not resolved merely by calling the subset “locally uniform.” Section II.259 only transports the theorem without loss on an actual finite subgroup or its coset.

V.213.3 Proposed proof-assistant modules

A formalization can be separated into five small interfaces rather than starting from the entire Gowers argument.

Finite weighted configurations. Define $\mathcal{Q}(w)$, $\mathcal{D}_w(\phi)$, and weighted distance as normalized finite sums, with an explicit proof that $\alpha > 0$ implies $\mathcal{Q}(w) > 0$. Represent a partial map by a total map together with the support of w . Establish invariance under subtracting an affine map and the exact graph energy dictionary for indicators.

Domain moments. Prove the autocorrelation identities by bijections of finite products and the triangle identity by the change of variables in Lemma II.254.2. Prove Fourier inversion, or reuse a compatible finite abelian Fourier API, for the pointwise link and marked-pair bounds. Check carefully that Fourier sums on the dual are not probability averages.

Finite modes. Only finitely many target values occur. A mode can therefore be chosen from a finite image without assuming that the whole target is a finite type. The zero-mass case must be defined explicitly. A reusable weighted collision-to-mode lemma then covers both the first and last selections.

Self-correction and the profile. Formalize the ordinary group-valued self-correction lemma separately. Its key statement is uniform in the fixed base point x ; an averaged-only version would be insufficient to prove additivity at every pair. The one-error inclusion–exclusion inequality is a five-case finite check. The Fourier marked-pair lemma must include both same-side and opposite-side placements in $a + b = c + d$.

The constant layer. Isolate (II.258.1)–(II.258.5) in a scalar lemma with $0 < \alpha \leq 1$, $0 \leq \lambda \leq 1/32$, and $0 \leq \varepsilon \leq 1/1024$. This prevents numerical bookkeeping from obscuring the structural proof. The exact rational inequalities in the supplied verifier indicate which parts can be discharged by arithmetic normalization.

Suggested new theorem names are

```
weighted_relative_affine_repair
relative_affine_defect_profile
quadratic_level_affine_repair
```

They are proposals, not declarations already present in the repository. The appropriate initial location is a new research subdirectory, with no change to the formalization-status ledger until compiled proofs are supplied.

[write] **7 October 2026; sources 54–59.** *Checked at HEAD fe7165a3c.* None of the three names is a declaration of the Lean development, and the ledger is unchanged by this source. The suggested subdirectory was not used: by its placement (7dc1ee29d) this source is source 56 of this report, with labels `gsr:fm:56:`, `gsr:en:56:` or `gsr:fz:56:` and files under the prefix `56-relative-affine-repair-`. `lemma_10_9` is at line 335 of `Section10.lean` and `lemma_10_9_holds` at line 717 of `Proofs10Shift.lean`; `proposition_6_1`, `lemma_6_2` and `corollary_7_6` at lines 37, 47 and 197 of `Sections06_07.lean`.

V.214 Further research questions

Source 56, delivered Section 12.

The results above leave several concrete problems. They are posed here as research directions; their inclusion is not a claim that each is a recognized open problem with an established name in the literature.

Question V.214.1 (The optimal density exponent). *What is the weakest U^2 hypothesis, as a function of α , that forces error-independent linear affine repair? Can the sufficient condition $u^4 \ll \alpha^5$ be replaced by $u^4 \ll \alpha^\theta$ for some $\theta < 5$?*

The pointwise triple-link bound explains the present exponent, but may be stronger than necessary. A likely route is to replace a worst-case link estimate by a link estimate averaged specifically over the error set, while still retaining an error term proportional to its mass. A lower-bound construction must fail repair, not merely saturate one Fourier inequality.

Question V.214.2 (Optimal fixed-uniformity soundness). *For fixed small $\lambda > 0$, determine the best coefficient in $t \leq C(\lambda)\varepsilon + o(\varepsilon)$, uniformly over densities, finite abelian groups, weights, and targets. Is the pointwise link imbalance in (F.61.8) the correct first-order obstruction?*

The present bounds place $C(\lambda)$ near $1/4$, but do not identify it. An extremal analysis would have to coordinate the error support with the phase of the triple convolution. The separate Fourier and inclusion–exclusion estimates may not be simultaneously sharp.

Question V.214.3 (A sharp relative cubic profile). *Can the coefficient 12 of t^2 in (F.61.8) be replaced by an exact target-torsion-sensitive expression, analogous to sharp full-domain profiles, with controlled Fourier-uniformity corrections?*

The current argument deliberately treats the six pairs uniformly and ignores cancellation patterns among two, three, and four erroneous labels. A refined count should distinguish positions on the same side from positions on opposite sides and should retain the additive energy of the marked error weight.

Question V.214.4 (The largest universal repair range). *How far can the numerical conditions $\lambda \leq 1/32$ and $\varepsilon \leq 1/1024$ be enlarged while retaining an explicit linear conclusion?*

The square-root expression in Proposition II.256.3 is better than the coarse estimate used in the main theorem. Optimizing it together with the local profile is a finite-dimensional problem. Such an optimization would improve constants but would not, by itself, constitute a new structural theorem. It should be accompanied by certified inequalities and counterexamples near any proposed boundary.

Question V.214.5 (Quadratic-domain rank). *Can the sufficient rank $R \geq 4r + O(1)$ in Theorem II.260.2 be lowered? Can exact Fourier distributions of specific quadratic systems replace the worst-case matrix-rank certificate?*

The factor four comes from the elementary estimate $u^4 \leq \sigma^2 \alpha$ combined with $\alpha \asymp p^{-r}$. A sharper calculation for individual systems may lower it substantially. One should distinguish improving a sufficient certificate from finding the true extension threshold.

Question V.214.6 (Higher-degree repair with a fixed domain condition). *Is there an analogue of the present weighted theorem for approximate polynomials of degree two or higher, with an explicit uniformity condition independent of the test error and a sharp first-order soundness constant?*

The natural first-order coefficient for a nondegenerate d -cube test would be 2^d , because that is the number of coordinate incidences. Turning that heuristic into a theorem requires a coarse higher-degree reconstruction and control of multiply marked cube configurations. Repeated vertices and low-characteristic polynomial phenomena must be handled explicitly.

Question V.214.7 (Propagation through a complete inverse proof). *What quantitative improvement survives if relative affine repair is used at every eligible high-agreement stage of a full inverse-theorem argument?*

The local gain can be substantial, but a global assertion requires tracking the rank, density, error, and sampling distributions at each application. In particular, inserting the theorem into a quadratic-repair proof does not automatically eliminate its other error-dependent rank requirements.

Question V.214.8 (Local domains without group structure). *Can a corresponding theorem be proved on regular Bohr sets or long progressions with explicit boundary errors and a clearly defined relative quadruple distribution?*

The exact triangle identity and the uniform self-correction translations are the first obstacles. A satisfactory statement should specify the regularity or approximate-invariance parameters and show how they enter the final linear error, rather than burying them in an undefined local-uniformity condition.

Question V.214.9 (Faster decoding and sampling). *Can the quadratic-time decoder be made nearly linear, or sublinear in a suitable oracle model, without introducing density-dependent label-query complexity? Can a weighted quadruple sampler be constructed efficiently from an explicit Fourier-uniform domain?*

The existence proof needs only modes with certified margins, suggesting sampling-based approximations. Uniform control over all ambient points and the cost of generating admissible pairs or quadruples are separate problems. A running-time statement must account for both.

Question V.214.10 (Minimal domain certificates and formal verification). *Which combination of autocorrelation variance, triangle variance, marked-pair energy bounds, and affine separation is sufficient when the domain does not have small U^2 norm? Can this abstract certificate be formalized once and instantiated for several distinct classes of domains?*


Proposition II.256.3 already uses the actual triangle variance. An abstract version might cover domains that fail the simple Fourier criterion but still have enough reliable local-to-global transport. The five-module formalization plan in Section V.213 is a concrete first step.

V.215 Proof audit and dependency ledger

Source 56, delivered Appendix A.

V.215.1 The dependency chain

The main existence proof has the following logical order:



Uniqueness is a separate consequence of Fourier coset separation. The quadratic-domain theorem uses only the main theorem, the spectral certificate, and the elementary quadratic Fourier bound. The property-testing corollary is the contrapositive of the main theorem in its small-error branch.

V.215.2 Points checked explicitly

No lost normalization. $\mathcal{Q}(w)$ and $\mathcal{D}_w(\phi)$ are normalized by N^3 ; B by N^2 ; r and T by N . Thus extending a bad pair to a triangle introduces α , not αN . Formula (V.213.2) is the conversion to the $|A|^3$ convention in the graph-energy step.

No assumption of uniform fibers. The distributions ν_h are normalized by their actual masses $r(h)$. Only the moment identities control the transition to the ambient measure. Zero fibers are assigned explicit values and never divided by zero.

No target Fourier transform. All Fourier analysis is on G . The target can be any abelian group. The proof uses addition, subtraction, equality, and modes of finite collections of target values.

No hidden odd-order hypothesis. The main theorem never divides a group element by two. Same-side marked-pair terms are bounded as complex Fourier sums, not incorrectly identified with opposite-side nonnegative sums. Only the quadratic matrix-rank corollary assumes odd characteristic, precisely in Lemma II.260.1.

No limit at zero error. The coarse theorem provides an approximation at $\varepsilon = 0$, and the positive local profile forces its error to vanish. Exact extension is therefore proved directly under the fixed domain hypothesis.

No circular use of uniqueness. The decoder produces an affine map before the separation lemma is invoked. Uniqueness does not enter the construction or the bootstrap.

No formal-verification claim. The finite program checks exact consequences on specified inputs. It does not quantify over all groups or all real weights. The universal statements are proved in the text and remain candidates for external review and subsequent formalization.

V.215.3 Prior versus present ingredients

The ordinary homomorphism self-correction method and the existence of linear repair theorems belong to prior work [49, 228]. Basic Fourier identities, quadratic Gauss bounds, collision-to-mode estimates, and inclusion–exclusion are standard tools. The proposed extension is their particular quantitative assembly into Theorem F.61.1, with weighted relative normalization, error-independent uniformity, the explicit profile, and the consequences stated here. The audit does not certify worldwide novelty of any individual intermediate identity.

V.215.4 Package contents and source control

The package contains the master `article.tex`, its included section files, the compiled `article.pdf`, a Makefile, a README, a source and integration note, a source-manifest JSON file, the reference decoder, the finite verifier, and the recorded verification output. It does not include copies of external papers or font files.

[write] 7 October 2026; sources 54–59. As shipped here: the main file and its section files and the PDF are not staged (their text is printed in this report); the Makefile, verifier, decoder, source manifest and outputs are shipped under the prefix `56-relative-affine-repair-` (README of this directory). The build record `data/56-relative-affine-repair-build_validation.json` names a single-file `gowers_relative_affine_repair.tex` “delivered separately”; it is not in the archive.

The repository is changing rapidly. The source manifest distinguishes a pinned GitHub snapshot from separately observed file blobs and records the arXiv versions inspected. The suggested new research-directory name avoids renumbering any existing source in the consolidated report. Integration into a formal proof ledger should wait for actual compiled formal proofs.

V.216 Further research questions and a proposed programme

Source 57, delivered Section 7.

The following questions are suggested by the proofs. They distinguish finite classification tasks from improvements requiring a new analytic argument. They are proposed research directions, not assertions of unresolved status throughout the literature. In particular, the general qualitative inverse and near-extremizer principles have extensive earlier theory; the targets below concern precise quantitative, structural, or formal interfaces.

V.216.1 The upper energy spectrum

Research question V.216.1 (The optimal Schur threshold). Let $N(s)$ be the least eventual threshold for the partition classification in Theorem II.292.3. Determine $N(s)$ and all exceptional positive real sets below it. Is the proved upper bound $N(s) \leq 5s - 1$ asymptotically sharp?

[write] 7 October 2026; sources 60–69. *Advanced by source 66: $N(s) \leq 3s + 1$ for every s and*

$N(s) \geq 2s + 2$ (its Theorem II.302.1 and Proposition II.302.4); its Conjecture V.249.1 is $N(s) = 2s + 2$ (item II.53).

The first exact values are $N(1) = 4$ and $N(2) = 6$. Small-cardinality exceptions can be continuous families: $\{1, t, 1 + t\}$ has deficit one for every $t > 1$, $t \neq 2$. Thus integer enumeration alone cannot settle the exceptional classification. A useful next proof should sharpen the argument that separates the first imperfect row from the final perfect rows, while keeping track of the *sizes* of their defects.

Research question V.216.2 (The fourth and subsequent energy levels). For each fixed r , determine the r largest distinct additive energies of an m -element torsion-free set for all sufficiently large m , classify every equality case, and determine the smallest valid cardinality threshold.

[write] **7 October 2026; sources 60–69.** *The fourth level is found by source 66: $M_m - 8m + 20$ for $m \geq 11$, attained exactly by two oriented templates (its Theorems II.308.1 and II.308.2); the fifth and later levels remain open (its Question V.249.5).*

The third-level proof reduces the equality analysis to endpoint costs zero, one, and two. The partition theorem supplies a finite list for every fixed endpoint cost. A remaining issue is to prove which costs can occur on a branch attaining a specified global energy level. The numerical ordering of candidate families must also be checked: different affine models may have the same energy polynomial.

Research question V.216.3 (Energy formulas for partition models). Develop a uniform formula for the eventual energy of $\{0\} \cup \{j + \lambda_j : 1 \leq j \leq n\}$ in terms of a fixed partition (λ_j) , and classify coincidences between the resulting energy functions.

[write] **7 October 2026; sources 60–69.** *Answered by source 66 for boundary-hole models: $D(A_H(N)) = r(m - 1) - \kappa(H)$ exactly (its Theorem II.304.1); the generating function of κ is its Question V.249.4.*

The partition theorem reduces this calculation to finitely many terminal patterns at each fixed deficit. Such a formula could turn the endpoint induction into a symbolic enumeration procedure. A satisfactory implementation should return both the candidate models and a proof of the range of n for which the ordering of their energies has stabilized.

Research question V.216.4 (A larger rigidity window). For each fixed K , does $D(A) = (M_m - E(A))/4 \leq Km$ force all but $O_K(1)$ points of a torsion-free m -set into an arithmetic progression of length $m + O_K(1)$? Determine the necessary exceptions and the best dependence on K .

Our bounded-Schur theorem controls an endpoint once its cost is bounded. A global bound $D(A) \leq Km$ controls a sum of endpoint costs; it does not immediately provide a long prefix with a bounded-cost terminal extension. Locating a useful prefix or suffix is therefore a distinct step. A complementary problem is the regime $D(A) = o(m^2)$, where a sharp proportion-of-edits statement may be more natural than a bounded number of edits.

Research question V.216.5 (Rectification and weights). Find explicit finite cyclic hypotheses under which the new equality classifications transfer through a Freiman isomorphism of order two. Can an ordered, nonnegative weighted defect identity support analogous stability conclusions for functions?

The cyclic problem must preserve both directions of every two-sum relation and justify recovery of an actual common step. Merely choosing representatives in an interval does not do so. The weighted problem would connect exact set classifications more directly to the bounded functions used in Fourier arguments; both the right extremal normalization and the equality family require care.

V.216.2 Quantitative extraction beyond Section 13

Research question V.216.6 (Propagating Fejér selection through higher degree). Can the general arrangement-selection theorem in source 40 be combined with the remaining Section 16 proper-box induction to produce a fully quantified improvement in Corollary 16.11?

The first concrete test is the next multidimensional extraction case, with every side length an integer, all boxes proper, and every density threshold uniform in the actual selected set. The present Section 13 proof shows why improving purification can matter substantially after a later spectral substitution. It does not supply the missing higher-dimensional geometric induction.

Research question V.216.7 (A better prepared frequency graph). Improve the exponent $d = 2\,996\,201\,429$ in Equation (II.290.2), while retaining both separate order-eight Freiman structure and an almost-everywhere arrangement identity.

The two successive fibre restrictions have a large cost before Fejér purification begins. One possibility is to exploit the shared origin of the rows and columns as derivatives of the same function. Another is to retain several large Fourier frequencies before making a coherent choice. As a concrete benchmark, a prepared density at least $a^{2^{30}}$ would lower the final power-of-two density exponent again with the current geometric interface. This is a target for a new argument, not a consequence of the existing constant rounding.

Research question V.216.8 (A fully explicit size threshold). Give a closed quantitative upper bound for $N_{\text{sq}}(\delta)$ and hence for the complete $N_0(\alpha)$ in Remark II.290.4.

The task includes the scale recurrence, integer cutoffs, properness, and the common-step extraction. The explicit bound $p \geq a^{-2^{37}}$ already proved here controls purification only. A successful result should identify which stage dominates the threshold and should avoid a circular dependence on the actual input density.

Research question V.216.9 (Sharper correction constants). Determine the best quantitative stability bounds for symmetric approximately additive tables and the consequent near-maximal U^4 estimates.

The proof uses a coarse correction on rows with a large test error, then a collision argument on slopes. Both are visible places to improve the constant 332. The cubic integration constant 6194 comes from a product of two average-deficit losses; a simultaneous integration argument may avoid that product. The perturbation in Proposition II.247.12 supplies a lower benchmark of 1/8 for the ratio of squared cubic-phase error to the Q_4 deficit, so the linear order is settled but its best coefficient is not.

Research question V.216.10 (Dependence on degree and characteristic). Improve the degree dependence in the all-degree phase approximation of Corollary II.247.11. Give comparably explicit estimates in small characteristic and on higher-dimensional vector spaces.

Our scalar-frequency correction costs only $O(s)$ at one step, but iterating degree descent multiplies those costs and gives $K_k \leq 2 \cdot 109^{k-2}(k-2)!$. Is a substantially smaller growth rate available by correcting several degrees together? For $p < k$, the classical polynomial demodulation divides by a vanishing factorial. The appropriate target must allow the relevant nonclassical polynomial phases. On $\mathbb{F}_p^n$, compatibility involving the base variable, not just symmetry of the derivative parameters, must be addressed.

Research question V.216.11 (From local extraction to global density increments). After the higher-degree geometric steps have been completed, what quantitative improvement survives a complete density-increment iteration with all parameter choices and thresholds propagated?

The desired output is one explicitly stated bound for $r_k(N)$, with a complete proof from the revised local estimates. Such a calculation must be compared with contemporary inverse-theorem methods, including [2], and cannot be inferred from a single improved exponent in Theorem 13.12.

V.216.3 Formalization and reproducible classification

Research question V.216.12 (A reusable formal interface). Formalize the purification-to-square synthesis and the independent near-maximal endpoint, then discharge the pinned `theorem_13_12` type without weakening its exponent or properness requirements. In parallel, formalize the reflection-row lemma and the partition classification.

These directions share little proof infrastructure and can proceed independently. On the Fourier side, the best reusable statements are phase normalization, finite-group additive correction, and a dominant-character selector with an average error bound. On the structural side, the natural core is an ordered finite set together with its nonnegative row defects. Exact finite enumeration can help test proposed statements, but should remain outside the logical dependency of the general proofs.

A practical order of work is: first formalize the proved Section 13 synthesis and the partition theorem; next settle the smallest new Schur thresholds and the fourth energy level; then attack higher-dimensional extraction and the complete quantitative iteration. This ordering exposes usable intermediate theorems before undertaking the longest analytic induction.

V.217 Proof dependencies, verification, and integration

Source 57, delivered Section 8.

V.217.1 A ledger of mathematical dependencies

The new results have three independent foundations. The following ledger distinguishes a cited substantial input from a proof reproduced within this article.

Component	Provenance	Treatment here
Fourier quadruples, order-eight restriction, arrangement abundance	Gowers, Proposition 6.1, Corollary 7.6, Lemma 12.4	Imported with exact constants
Fejér purification	Repository source 40	Specialized complete proof
Proper common-step square	Pinned uniform Section 13 extraction	Exact imported interface
Improved preparation and exponent propagation	This article	Complete proof
Near-maximal correction and phase integration	Classical correction background; explicit bounds here	Complete elementary proofs
Energy recurrence and preceding gaps	Hegvári; repository source 45	Required proofs reproduced
Bounded Schur classification and third-level equality	Additions relative to source 45	Complete proofs and sharp small cases

The principal inspected formal files are relative to the directory

`Combinatorics/Ramsey/Lean/GowersSzemerédi`

and comprise:

- `Definitions.lean` and `Sections12_13.lean`, for uniformity, derivative conventions, the context, and the exact catalogue type;
- `Proofs13UniformSquareExtraction.lean`, for the density-uniform extraction theorem;
- `Proofs13CompleteSquareExtraction.lean`, for $\gamma(\delta)$ and the constructed integer geometry;
- `Proofs13ExplicitFourierExponent.lean`, for the earlier explicit bound and its prefactor estimate;
- `Proofs13PrintedExponentComparison.lean`, for the strict comparison showing why the earlier density substitution does not establish the printed exponent.

The two principal input blobs are:

- Complete-square: `9e461b7be2e833f0c8e5f9899aba19360f611e61`.
- Uniform-square: `31d4c858f6761dca4ff71b2770e35f9863cc7c7e`.

The accompanying provenance file records the remaining identifiers.

The previous strict exponent comparison remains correct for its previous input density. Our proof changes that input and adds an independent endpoint, so it is consistent with the comparison theorem. The common-square input requires a sufficiently large prime depending on a *lower* density bound; the passage through ρ_0 and δ preserves exactly that quantifier order.

V.217.2 What the verification checks

The package contains two standard-library Python programs and their recorded JSON outputs. All assertions use integers, rational numbers, or exact finite combinatorics. No floating-point comparison of towers of exponentials is used.

1. `verification/quantitative_exact_checks.py` checks the fibre exponents, the purification count and threshold, the ceiling in d , the final 2^{41} and 2^{53} bounds, the printed-exponent overlap, and the endpoint constants. It also checks twelve finite Fejér parameter choices, 1 249 finite maps for the additive correction calculation, and 636 slope maps for the bilinear collision inequality. Among the additive maps tested, 36 nonzero-defect examples satisfy the strict correction premise. Further checks cover the cubic and quadratic polynomial identities, all 120 pairs of distinct four-cube vertices, and the finite-group correction recurrence. The all-degree checks include K_k through $k = 12$, coefficient identities through derivative order five, and cube-pair rank certificates for orders two through six.
2. `verification/verify_structural.py` examines all n -subsets of $\{1, \dots, 20\}$ for $5 \leq n \leq 9$, a total of 425 714 sets, for the deficit-two classification. It checks 257 754 sets containing zero inside $\{0, \dots, 20\}$, of sizes six through nine, for the third energy level, modulo dilation and reflection. It additionally constructs all 66 partition models for deficits one through eight at the theorem’s respective thresholds.

Every recorded check passed. The bounded searches do not establish classification over the reals or for arbitrary cardinality; those claims follow from the reflection and endpoint proofs. Likewise, the finite additive tests do not establish the correction lemmas; the majority and collision arguments do.

Separate AI-assisted proof passes examined the quantitative propagation, the repeated-vertex selection issue, the endpoint normalization and cocycle, the all-degree correction, the small Schur exceptions, and all branches of the third-energy equality proof. The enclosed review record describes those checks. This is an internal mathematical audit, not external peer review or a new machine-checked proof.

V.217.3 Suggested repository integration

The main TeX file is standalone and has no dependency on the copied source papers or on external figure assets. Its theorem labels use the prefixes `main:`, `qt:`, `ep:`, `st:`, `rq:`, and `audit:`. These can be renamed mechanically if the article is merged into an existing combined manuscript.

A written-proof intake can record Theorem F.62.1 as a route to the exact pinned Theorem 13.12 type with stronger density. A formal ledger entry should change only after the corresponding Lean proof is compiled and its dependencies audited. The two questions in source 45 can receive dated pointers to Theorems II.292.3, II.292.4 and II.292.7, while retaining their original wording as provenance.

For future Lean work, useful module boundaries are:

1. finite-group phase normalization and additive correction;
2. symmetric frequency selection and the near-maximal endpoint;
3. the specialized Fejér selection, followed by exact exponent arithmetic and the common-square synthesis;
4. ordered reflection rows, partition classification, and endpoint transfer;
5. the third-energy equality theorem and its six-point exception.

The present package supplies the mathematical statements and proofs for those modules. It does not include placeholder Lean declarations or alter the repository's current proof status.

[write] **7 October 2026; sources 54–59.** What became of these suggestions (checked at HEAD fe7165a3c). `FORMALIZATION_STATUS.txt` cites Theorem F.62.1 as the route to the printed Theorem 13.12, and the repository then proved `theorem_13_12` formally along it (note after Theorem II.290.3), so the first three module groups listed above have formal counterparts in the repository's own form; the partition and third-energy modules have none. The two questions of source 45 received dated pointers in this report (after its Question V.166.1). The labels were renamed with the prefixes `gsr:fm:57:`, `gsr:en:57:` and `gsr:fz:57:` in front of the delivered ones.

V.218 Examples, numerical exploration, and exact checks

Source 58, delivered Section 12.

V.218.1 Why both torsion and repetitions need explicit treatment

Example V.218.1 (A repeated bad tuple). Let $G = \mathbb{Z}/2\mathbb{Z}$, $H = \mathbb{Z}/4\mathbb{Z}$, $B = \{0, 1\}$, and $\phi(0) = 0$, $\phi(1) = 1$. The tuple $(0, 0, 1, 1)$ is additive in G , but its target discrepancy is $-2 = 2 \neq 0$ in H . For one layer with $K(t) = \cos^2(\pi t)$, its actual containment probability is $1/4$. The formal occurrence product instead has mean $9/64$.

Indeed, averaging over the target character makes the nonzero frequencies of the second distinct-point factor vanish, so the actual average is $c_0^2 = 1/4$. For the occurrence product the two factors are each squared. Their Fourier supports lie in $[-2, 2]$, and averaging the target character again leaves only frequency zero for the second factor; its constant coefficient is $R_2 = 3/8$. The remaining shift average is another R_2 , giving $9/64$. The lower occurrence-product value cannot be used as an upper bound on bad survival.

Example V.218.2 (A noninvertible additional relation). In $\mathbb{Z}/N\mathbb{Z}$ with $2 \mid N$, a nonprincipal relation can leave an equation $2x = a$ after the other variables are fixed. Its nonempty fibres have two elements. Over a prime field of odd characteristic, the same completion would be unique. Lemma II.280.1 retains this fibre size through $|G[2]|$, rather than silently importing field rank counting into a composite modulus.

Example V.218.3 (Degree-one amplification has an exponential arity cost). For $K(t) = \cos^2(\pi t)$ the only nonzero coefficients are $c_0 = 1/2$ and $c_{\pm 1} = 1/4$. Thus

$$P = 2^{-m}(1 + 2^{1-m}), \quad Q = 2^{-m}, \quad \mathcal{R} = 1 + 2^{1-m}.$$

The principal exponent is

$$a_K = \frac{m \log 2 - \log(1 + 2^{1-m})}{\log(1 + 2^{1-m})} \sim (m \log 2) 2^{m-1}. \quad (\text{V.218.1})$$

Replacing the degree-one filter by a suitable powered cosine changes this exponential dependence into $O(m \log m)$. This calculation isolates the relevant mechanism; it does not purport to reconstruct every constant in Gowers's original proof.




Figure V.218.1: Numerical exploration of powered cosine filters. The left panel shows the dependence on the degree at $m = 16$. The right panel normalizes the best sampled exponent by $m \log m$. The dashed reference line is the limit for the infimum over the full filter class, $1/(2 \log 2)$, proved in Theorem B. It is not a claimed limit of the degree-restricted finite search. Numerical data select candidates; the exponent-54 conclusion uses exact integer certificates.

V.218.2 A reproducible finite search

For each tested arity, the exploration evaluates every degree $1 \leq q \leq 3m$. The tested arities are

$$m \in \{4, 8, 16, 32, 64, 128, 256\}.$$

The table gives the best value *within that tested range* for a_K with worst discrepancy order two. The reported decimals are rounded numerical values.

m	Best sampled q	a_K	z_K
4	5	7.755382	4.996779
8	12	21.177067	17.854033
16	25	52.206426	48.341189
32	53	122.606592	118.213699
64	111	279.974503	275.060702
128	231	627.704888	622.274215
256	475	1388.924965	1382.979512

The script uses logarithmic sums and evaluates $\log(P/Q)$ through the odd-frequency mass with `log1p`, avoiding cancellation when P/Q is close to one. Very small masses are stored in scientific notation derived from their logarithms. The generated CSV files contain all 1,524 sampled candidates, not only the selected rows.

V.218.3 What was checked computationally

The finite verification has two distinct purposes. First, `verify_constants.py` checks the five finite integer inequalities used in the proof of Theorem A, together with the final threshold and retention comparisons. These calculations are exact and are part of the explicit numerical certificate.

Second, `verify_expectations.py` directly computes small-group survival probabilities using one Fourier factor per distinct point. For $m = 4$ it covers cyclic domain orders 2, 3, 4, 5, 6, 8, 9, target orders 2, 3, 4, 6, cosine degrees 1, 2, 3, and layer counts 1, 2, 3. It uses full and punctured domain sets with quadratic, step, and spike maps. The 1,512 cases give 6,048 principal inequality checks; 1,188 cases contain repeated bad tuples, and 615 require a nonprincipal correction beyond the principal bad contribution. All pass in rational arithmetic. The repeated-tuple regression in Example V.218.1 is checked separately.

These small cases test implementation-sensitive identities and catch the incorrect per-occurrence treatment of repeated points. They are not an exhaustive verification of the general theorem. The latter is established by the character expansion, fibre bound, Hölder estimate, and signed objective proved above. Likewise, the numerical search does not prove a global optimum over filter degree.

V.219 Further research questions and a proposed programme

Source 58, delivered Section 13.

The results settle a specific local extension question and establish an asymptotic barrier for a specified method. They also isolate several substantive problems that remain. The following questions distinguish finite optimization, structural improvements, and propagation through the larger argument.

Question V.219.1 (The sixteen-variable optimum). *Determine or tightly bound $I_{16,2}$, including optimization beyond powered cosine filters. Can one certify an admissible filter with $\mathcal{A}_{16,2} < 52$, and then improve the clean integer exponent after auditing its finite-size constants?*

[write] **7 October 2026; sources 60–69.** *The certificate is supplied by source 68: an admissible degree-twelve rational filter of principal cost $51.445 < 52$ with exact integer certificates, which lowers the exponent to 53 (its Proposition II.312.1 and Theorem II.312.2); $I_{16,2}$ itself is still undetermined (its Question V.256.7).*

The current value $52.2064\dots$ is a candidate value, not a lower bound. A useful computational project would parameterize nonnegative trigonometric polynomials through a sum-of-squares representation, optimize the nonlinear Fourier objective, and turn the final candidate into rational positivity and integer certificates. Checking nonnegativity only on a numerical grid would be insufficient.

Question V.219.2 (The next asymptotic term). *Find the first correction to the leading asymptotic for $I_{m,a}$, and identify the scale of the degree and the shape of near-optimal filters.*

[write] **7 October 2026; sources 60–69.** *Advanced by source 68: uniform two-parameter bounds, with the fixed- a upper error $O_a(m(1 + \log \log m))$ and the growing-order limit $1 + 1/(2c)$ (its Theorem II.313.1); whether the error is $O_a(m)$ is its Question V.256.8.*

The choice $q = \lceil m(\log m)^2 \rceil$ is designed for a short proof of the leading term. It is not a second-order optimizer. Sharper control of the residue-class imbalance of powers of binomial

coefficients, combined with a stronger coefficient concentration inequality, is a concrete route to reducing the gap between the two bounds. The fixed- a theorem also leaves the regime in which a grows with m open.

Question V.219.3 (Beyond independent positive filters). *Can restriction on every cyclic group have a retention exponent $O(m)$, uniformly over arbitrary abelian targets, by a method outside the independent positive Fourier-filter class?*

Theorem B rules out a linear principal-spectrum cost for the certificate analyzed here, including heterogeneous independent scalar layers. It does not rule out better use of nonprincipal terms or input structure, adaptive selection, joint filters on several characters, signed Fourier constructions with controlled cancellation, or a structural decomposition before restriction. The prime-field results indicate that linear dependence is a meaningful benchmark, but do not prove that it is achievable uniformly in small torsion.

Question V.219.4 (Sharp finite-size thresholds). *What is the best sufficient point-count threshold for a prescribed retention coefficient? Can the collision profile and the coefficient-weighted torsion profile substantially improve (II.282.13), especially at $m = 16$?*

The exact certificate (II.282.4) already avoids several unnecessary losses, and (II.281.2) makes an improvement dependent on the actual tuple geometry. A complete answer should also give examples showing which losses are necessary. Small sets can always retain a respected singleton, but this alone does not bridge the entire intermediate-size range of the present theorem.

Question V.219.5 (Low-exponent domains). *Find a useful analogue for families such as $(\mathbb{Z}/2\mathbb{Z})^s$ in which the domain torsion factor in (II.282.3) is comparable with $|G|$.*

A field-rank argument naturally belongs to these families, while the integer-minor bound used here is suited to cyclic groups and bounded invariant-factor rank. A hybrid theorem that detects the effective rank of the coefficient system in each primary component could join the two regimes. Merely substituting the ambient Smith kernel count into a bound for B^m would not be justified.

Question V.219.6 (Systems of relations and higher arrangements). *Extend the decaying multi-layer error to a general integer relation matrix while selecting one common subset and counting every tuple, including all collision patterns.*

Shkredov's system-level theorem and the repository's finite-field arrangement results are important starting points. Over composite groups, one must track saturated relation lattices, ranks of added Fourier rows, Smith invariants, and target discrepancy subgroups. The single-minor proof here settles only the case of one balanced relation. A useful first target is two independent balanced relations with a fully explicit error in terms of the relevant minors.

Question V.219.7 (Using target-order information). *Given the distribution u_a of discrepancy orders, optimize the certificate based on $\sum_a P_a^k u_a$. Can several filters be chosen to address different parts of this distribution with a smaller finite-size cost?*

The worst-order-two bound discards information when most bad discrepancies have large order. The mixed-filter formula is already available, but optimization must account for both the principal gains and the moment-dependent errors. Theorem B bounds only the principal cost for a fixed possible order, not this full data-dependent problem.

Question V.219.8 (Efficient certified construction). *Can one construct a qualifying subset and an a posteriori certificate in time substantially below enumerating all b^m tuples?*

The proof is finitely constructive by exhaustive conditional expectation, but it gives no efficient runtime. A practical derandomization should estimate the signed objective, exploit convolution or sparse graph structure, and retain rigorous error bounds for the unexamined tuple mass. Proving a runtime bound is a separate requirement from proving that a finite search terminates.

Question V.219.9 (End-to-end quantitative propagation). *What improvement survives after the new restriction constants are carried through the later stages of Gowers’s argument, with every size threshold and parameter dependence recalculated?*

This is the necessary next step before claiming an improved bound for progression-free sets from this route. The retained configuration mass, retained point density, accuracy parameter, and starting modulus must be propagated together. A smaller local exponent alone does not determine the final bound, and it need not compete with modern global arguments.

Question V.219.10 (Formal certificates and reusable interfaces). *Can the expectation theorem be formalized once for arbitrary finite abelian groups and admissible filters, with the numerical corollaries generated by exact certificates?*

This would separate the reusable mathematics from parameter selection. The most delicate reusable lemma concerns actual point selection for repeated tuples; the most valuable arithmetic interface concerns coefficient-weighted kernels of multiplication maps. The five certificates at $m = 16$ supply a small first instantiation.

Suggested order of work. First formalize and integrate the central expectation theorem and Theorem A. In parallel, pursue certified finite filter optimization and the collision-aware threshold, since those projects have explicit numerical targets. Next develop a two-relation composite-group case and an algorithmic certificate. Finally, propagate the best validated local statement through the later proof before assessing any global quantitative claim.

V.220 The essential integer certificate

Source 58, delivered Appendix A.

The following complete standard-library calculation verifies all finite numerical inequalities used in the proof of Theorem A. The symbolic uses of $0 < t \leq 1$ and $0 < \eta \leq 1$ are proved in the text, rather than inferred by sampling those parameters.

```
from math import comb
from fractions import Fraction

A = sum(comb(50, 25+n)**16 for n in range(-25, 26))
B = sum(comb(50, 25+n)**16
        for n in range(-25, 26) if n % 2 == 0)

assert A >= 2**750
assert 250*A >= 481*B
```



```

assert 481**53 >= 2**50 * 250**53
assert 481**50 >= 2**47 * 250**50
assert comb(100, 50) <= 2**97

# Domain torsion plus repeated-position pairs:
assert 4 * (50 + comb(16, 2)) == 680

# Ratio of the sufficient size bound to (4/t)**54:
# (680 * 2**97 * t**(-51)) / (2**108 * t**(-54))
# = (85/256) * t**3 <= 85/256 < 1.
assert Fraction(680 * 2**97, 2**108) == Fraction(85, 256)
assert Fraction(85, 256) < 1

# Certified retention divided by (alpha*eta/4)**54
# is 8/eta, hence at least 8 when 0 < eta <= 1.
assert Fraction(2**108, 2**105) == 8
print("All exact certificates passed.")

```

The companion `data/exact_constants.json` records the complete integer numerators as decimal strings, together with every check and its interpretation. No stochastic numerical computation enters this certificate.

V.221 Package contents and reproduction

Source 58, delivered Appendix B.

The archive contains this article’s L^AT_EX source and compiled PDF, the source audit, exact verification programs and their output, numerical exploration data, the figure and its generator, and build instructions. Run the following from the extracted package directory:

```

python3 code/verify_constants.py
python3 code/verify_expectations.py
python3 code/explore_kernels.py --no-figure
latexmk -pdf -interaction=nonstopmode -halt-on-error \
    gowers_torsion_uniform_restriction.tex

```

The verification programs and numerical exploration without plotting use only the Python standard library. Regenerating the figure requires Matplotlib and its NumPy dependency; the supplied vector PDF is already sufficient for compiling the article. A conventional TeX Live installation with the packages named in the preamble is sufficient.

The report was prepared with OpenAI ChatGPT, including independent parallel proof reviews and exact computational checks. The final mathematical statements remain subject to ordinary external review; the package makes no claim of peer review or Lean certification. Its principal source comparison is pinned to the repository commit specified in Section F.63. No remote repository content is modified by this deliverable.

V.222 Reproducibility and verification scope

Source 59, delivered Section 11.

The accompanying archive contains a standalone L^AT_EX source, the compiled PDF, three standard-library Python programs, their exact JSON reports, an input-source manifest, build instructions, and a detailed test inventory. The source embeds its bibliography and needs no external figures or private style file. The verification programs are supplementary checks; the general theorems rest on the proofs above.

The principal extraction verifier performs the following exact checks:

- 149,998 universal/odd one-element probability inequalities, 79,964 odd-order cutoff cases, and 72 prime-primary cases; adaptive recurrences for the specified ranges in the inventory.
- 8,750 separator applications to 8,508 distinct coordinate-model/forbidden-set pairs. Different theorem variants sometimes revisit the same input.
- All 2,441 partial-map cases in six small-domain suites, yielding 6,495 applications of universal, odd, or prime-primary extraction. For every selected subset it directly checks all Freiman orders up to the requested order, including repeated summands.
- 16,459 character tuples in projective rank-obstruction checks, 16 exact Laurent constant coefficients, and 18 finite-torsion Riesz products with their 36 deleted products.

The affine verifier evaluates 23,516 map cases: 21,116 exhaustive cases in explicit cyclic/product models, plus 2,400 reproducible sampling evaluations. Samples can repeat maps; this is not a count of distinct isomorphism classes or distinct functions across all suites. It computes graph energy and signed-convolution energy separately, verifies the exact expansion and all stated quartic bounds including the support gains, and checks the final local inequalities directly. The original-kernel equality classification is checked on all 23,489 evaluations with nonempty support; the endpoint checks cover 2 eligible four-ninths evaluations and 104 eligible no-order-three half-density evaluations. The equality checks use translated-subgroup closure for the support, independently of the energy criterion for a coset. It also checks 9,760 signed-label quadruples and the one-point identity on 350 specified maps with domain sizes 16 through 40.

The separate realization verifier checks small instances of the prescribed-relation construction against exhaustive multisets and subsets, including finite-domain reduction and the exact weighted optimum. Its JSON report records the exact instance list and counters, so these checks are not silently included in the preceding totals.

Important limits of the computations are explicit. They do not establish universal maximality by exhaustive search on all groups. They do not numerically verify infinite-target Haar arguments, arbitrary real weights, or the general finite-image rounding theorem. Those are analytic results. The one-point checks verify the claimed energy values, while maximality for $N \geq 16$ follows from the recovery theorem. The programs reject Python optimization flags that would disable assertions. A successful run is a consistency check, not a formal proof or a claim of mathematical priority.

For integration, retain the hypothesis $W[2] = \{0\}$ in the affine layer, keep the additional $W[3] = \{0\}$ assumption on the half-range theorem, preserve the separate endpoint conclusions, and count sign orbits correctly when involutions occur. The one-color/coset classification for the original averaged kernel must not be transferred to its strengthened support-defect bound, which has additional equality cases. The target's finite image size and its exponent are different parameters. These distinctions carry the mathematical content of the refinements.

V.223 Further research and a formalization programme

Source 59, delivered Section 12.

Except for the explicitly named lonely runner conjecture, the questions below are proposed here as follow-ups to the proved results. Some refine questions in the repository; no claim is made that their formulations have mathematical priority. The general mixed-torsion extraction problem from source 24 now has an answer independent of the target's exponent and rank in Theorem II.264.1. The support-coset stability question at scale η/t^3 was already answered in substance by combining sources 11 and 14, as the repository records. Those existence questions should therefore be replaced by sharper quantitative targets.

Question V.223.1 (Proposed: radius at the exact prime-primary rank). *For a prime p and $q \geq 2$, set $m_p(q) = \lfloor \log_p((p-1)q+1) \rfloor$. Determine*

$$R_p(q) = \inf_{H,F} \sup_{\chi_1, \dots, \chi_{m_p(q)}} \min_{f \in F} \max_j \|\chi_j(f)\|_{\mathbb{T}},$$

where H ranges over finite abelian p -groups and F over symmetric nonzero subsets with exactly q sign orbits. Theorem II.263.4 gives $R_p(q) \geq 1/(2p)$. Is this the optimal radius, and which forbidden sets attain or approach the infimum?

Sharpness of the Haar survival estimate does not answer this question: the optimizing characters may depend on the entire set. A small concrete test is $R_3(3)$. For $F = \pm\{1, 2, 3\} \subset C_9$, the minimum distances for character parameters $0, \dots, 8$ are $(0, 1, 2, 0, 1, 1, 0, 2, 1)/9$, giving the upper bound $2/9$.

Question V.223.2 (Proposed: optimize radius and character-image cost jointly). *Given a finitely generated abelian group H , an explicit forbidden set F , and $k \geq 2$, optimize the guaranteed retained fraction over separating families with unequal radii. For a character of image order N , its contribution is*

$$c_k(N, \theta) = \frac{\lceil N\theta/k \rceil}{N}$$

from Theorem II.264.3; for infinite image it is $c_k(\infty, \theta) = \theta/k$. Which order profiles or relations among defects allow the product of these factors to improve substantially on every choice of a common radius?

A useful solution would return a verifiable separating family and its actual localization cost. Minimizing the number of characters alone can miss the benefit of small finite images. The exact profile in Theorem II.263.1 supplies one certificate against which a more adaptive optimization can be compared.

A concrete first stage is to compare all relevant rational radii in small groups of exponent p^2 or p^3 , while recording character-image orders. The desired output is an exact certificate with its retained fraction and a proof of optimality within the chosen class of localizations. Numerical sampling alone cannot certify an optimal radius or distinguish an endpoint from a nearby failure.

Question V.223.3 (Proposed: sharp retention from a prescribed number of defects). *For fixed $k \geq 2$, let $\lambda_k(q)$ be the infimum over finite nonempty function graphs whose order- k vertical defects have at most q sign orbits, of the largest proportion of the domain supporting an exact Freiman restriction of order k . Is there a constant $c_k > 0$, independent of both groups, such that*

$$\lambda_k(q) \geq \frac{c_k}{q+1} \quad (q \geq 0)?$$

If not, determine the correct decay and construct graph examples attaining it.

The present universal lower bound is $(5k)^{-\lfloor \log_2(q+1) \rfloor}$. Theorem II.265.4 realizes every symmetric forbidden set as graph defects, with optimum retained proportion $1 - 1/(2k)$. Thus sharp character rank alone does not settle retained mass. A lower-bound example must control every exact restriction. The explicit realizations also have difference doubling comparable to their size, leaving small-doubling examples as a distinct quantitative challenge.

There is also a weighted version: replace cardinality by arbitrary nonnegative weights on the domain and maximize retained weight. The theorem here covers this version. Determine whether its optimum agrees with the unweighted optimum, or whether highly concentrated weights expose an additional obstruction that is invisible in uniform examples.

Question V.223.4 (Proposed: a linear dependence on the Freiman order). *For $k \geq 2$, define $a(k)$ as the infimum of exponents $a \geq 0$ for which a constant $c_k > 0$ guarantees an order- k restriction of relative size at least $c_k C^{-a}$ whenever $|\Gamma - \Gamma| \leq C|\Gamma|$. Can one prove $a(k) = O(k)$? Corollary II.265.2 gives*

$$a(k) \leq (2k + 1) \log_2(5k) = O(k \log k).$$

Alternatively, are there graph families forcing $a(k)/k$ to be unbounded along a sequence of orders?

This asks for the final retained size. One possible route combines defect counting with localization before either is bounded separately. It remains meaningful even if the general estimate $|F \cup \{0\}| \leq C^{2k+1}$ cannot be improved. A positive answer to Question V.223.3 would already give the desired linear upper bound.

An alternative approach should identify a usable invariant beyond the cardinality of the vertical fibre, such as its additive dimension or distribution of orders. Any improvement must state the relation between that invariant and C ; treating a favorable profile as automatic from small doubling would leave the central implication unproved.

Question V.223.5 (Established conjecture; proposed small-doubling restriction). *The classical lonely runner conjecture asserts that, for distinct positive integers $v_1, \dots, v_q$, some $t \in \mathbb{T}$ satisfies*

$$\min_i \|tv_i\|_{\mathbb{T}} \geq \frac{1}{q+1}.$$

For a concrete restricted target, let $k = 2$, let all nonzero vertical defects of Γ have infinite order, and suppose they have $q \geq 2$ sign orbits. If $|\Gamma - \Gamma| \leq q^{1/2}|\Gamma|$, must one character separate all these defects at radius $1/(q+1)$?

The named conjecture is an established open problem [239]; the displayed small-doubling restriction is proposed here. By Theorem II.265.4, merely requiring graph defects leaves every integer forbidden set available, so a substantive extra condition is necessary for this approach. Bedert's improved general estimate already transfers through Proposition II.266.2 [238]. Invoking that estimate must be distinguished from proving an improvement to it.

Question V.223.6 (Proposed refinement: the maximal local affine range). *Let $\ell : G \rightarrow W$ be a specified affine reference, let $t = |\{f \neq \ell\}|/|G|$, and write $\Psi(t) = 4t - 10t^2 + 6t^3$. Under $W[2] = \{0\}$, does $\varepsilon(f) \geq \Psi(t)$ continue throughout $0 < t < 1/2$? Under $W[2] = W[3] = \{0\}$, does it hold throughout the positive-cubic range $0 < t < 2/3$? In any larger valid interval, does equality still force one nonzero error value on a subgroup coset?*

Theorem II.250.4 proves the respective ranges $t < 4/9$ and $t < 1/2$. The question concerns a specified reference; assuming it is a nearest affine map would change the problem. The next step should retain the interaction between the signed four-error term, the three-error terms, and the support geometry, rather than optimizing their bounds independently.

Small-group searches can guide this question if they record both the specified reference and the true nearest-affine distance. A counterexample to one formulation may satisfy the other. Candidate improvements should therefore be accompanied by the exact signed-pattern inequality they require, and candidate obstructions by complete finite maps with exact energy counts.

Question V.223.7 (Proposed: the optimal support–color tradeoff). *For a fixed rational $0 < t < 1/3$, determine the set of pairs $(\alpha, \beta) \in [0, \infty)^2$ for which every map into a target with $W[2] = \{0\}$, at support density t from a specified affine reference, satisfies*

$$\varepsilon(f) - \Psi(t) \geq \alpha a + \beta \left(t^3 - \frac{E(S)}{N^3} \right).$$

Which pairs on its upper boundary are attained or approached by finite examples, and how does this boundary change if $W[3] = \{0\}$ as well?

The present bounds supply the pairs $(2 - 6t, 6)$ and $(2 - 9t/2, 3)$ in the first target class, together with all their convex combinations and coordinatewise decreases. Without order-three torsion, the enlarged-range argument supplies $(\min\{2 - 3t, 4 - 8t\}, 4)$. Optimizing this region asks how much support structure can be controlled simultaneously with color concentration; optimizing either coefficient in isolation can discard the interaction. The endpoint equality cases themselves are settled in Corollary II.250.6: density $4/9$ gives strict inequality in the first class, and density $1/2$ in the second class has precisely the one-value index-two-coset extremizers.

Question V.223.8 (Proposed refinement: involutions in the target). *When W contains a nonzero involution, determine the largest local range for the appropriate profile*

$$\Psi_{\text{inv}}(t) = 4t - 12t^2 + 8t^3.$$

Can the averaged signed-kernel method enlarge the repository’s sufficient range $t < 1/3$ while retaining explicit nonnegative remainders that distinguish unequal colors, opposite colors, and involution-valued colors?

For color multiplicities n_b , useful separate quantities are $A = s^2 - \sum_b n_b^2$ and $J = \sum_b n_b^2 - \sum_b n_b n_{-b}$, where $s = \sum_b n_b$. The existing involution theorem already gives a cubic with these remainders and classifies equality below $1/3$. The unresolved target is a larger range and sharper joint coefficients, including the interaction with order-three cancellations.

Question V.223.9 (Proposed refinement: the remaining small-domain energy maxima). *For $W \neq \{0\}$ with $W[2] = \{0\}$ and $2 \leq |G| < 16$, determine*

$$M(G, W) = \max_{f \text{ non-affine}} \mathcal{E}(f)$$

and classify the extremizers. In particular, does the one-point formula from Corollary II.251.3 hold at orders 13, 14, 15? Organize any exceptions through the homomorphisms and quotients of G , including failure to extend a homomorphism from a subgroup.

A cutoff cannot simply be removed: for $G = C_{12}$ and $W = C_3$, the map equal to 1 on odd inputs and 0 on even inputs has energy 1296, whereas the one-point formula gives 1266. The coset formula verifies the first count; the two-point image rules out affinity. This example already lies in the involution-free target setting of the manuscript.

Question V.223.10 (Proposed refinement: higher graph energies). *For $r \geq 3$, define*

$$E_r(f) = |\{(x_1, \dots, x_r, y_1, \dots, y_r) : \sum_i x_i = \sum_i y_i, \sum_i f(x_i) = \sum_i f(y_i)\}|.$$

Determine the sharp small- t envelope of $1 - E_r(f)/|G|^{2r-1}$ at distance t from a specified affine reference, starting with $r = 3$. Which support geometries and error-value patterns give equality, and precisely which torsion assumptions are needed?

The first investigation should keep repeated coordinates and relations among two or three error values explicit. Existing near-maximal set-energy and cube-collision stability theorems provide useful inputs, while this question requires the full sparse-error expansion for a function graph. Any application to Gowers's later respected arrangements must identify the actual counting tensor used there.

One useful intermediate deliverable is a complete three-sum expansion for errors supported on a coset, allowing two nonzero colors. This isolates the first competition between support structure and torsion in the values. It can then be compared with arbitrary supports before proposing a universal extremal formula or an equality classification.

V.223.1 A numbered priority path for formalization

The following is a proposed order of work. It asserts no existing Lean proof of the new results.

1. **Fix the finite interfaces.** Define graph defects, sign orbits including singleton involutions, error supports, color multiplicities, and ordered energies. Prove the basic bijections with cardinalities in $\mathbb{N}$ before normalizing in $\mathbb{R}$.
2. **Formalize exact separation for finite targets.** Prove the cyclic arc count, strict averaging with a distinguished positive-mass outcome, and the adaptive deletion recurrence. Express the rank bound first by integer inequalities $b^m \leq (b-1)q + 1 < b^{m+1}$; derive the logarithmic notation afterwards. Verify the projective-line sharpness construction independently.
3. **Connect separation to graph extraction.** Establish weighted localization, finite-image rounding, the vertical-fibre injection, and the precise sum-and-difference growth interface. Check explicitly that the auxiliary minimal-growth subset in the classical growth proof does not replace the graph in its conclusion. Extend to finitely generated targets after the finite version is complete.
4. **Verify the signed energy identities.** Formalize the exact expansion, the three-position kernel average, and the support and color remainders. Keep the hypotheses excluding orders two and three visible in the statements where they enter. Derive the new local ranges and their open-range equality classifications from these identities.
5. **Compose the recovery and stability results.** Supply or import verified statements for homomorphism correction and coset stability with their exact normalizations. Then prove the recovery thresholds, the size-16 corollary, and the near-extremizer composition. The endpoint

classifications proved here should be included; the optimal support–color frontier remains a separate research task.

6. **Audit the downstream uses.** Instantiate the proved interfaces at the actual uses in Sections 6 and 7, retaining a ledger of domains, losses, and thresholds. Only a completed audit of subsequent dependencies could support a new end-to-end bound. Exact computation scripts can serve as regression checks throughout; they do not supply kernel verification of the general statements.

V.224 Consequences, comparison, and integration

Source 60, delivered Section 10.

The results have different logical roles. The norm theorem is a uniform asymptotic comparison on varying groups. The arrangement theorem is a sharp local repair result on a full product. The two-direction theorem repairs the compatibility between two such full-product models. The restriction theorem applies to a subset and gives a quantitative improvement at the earlier approximate-homomorphism selection step. Keeping these roles separate prevents a local improvement from being assigned a conclusion that requires additional structural arguments.

V.224.1 A precise record of which questions are answered

Table V.224.1: Status relative to the inspected repository.
“Solved” refers to the stated local or limiting problem.

Prior target	Result here	Remaining qualification
Source 29: limit of c_d^{odd} ; could it be $1/2$?	Theorem III.99.1 identifies the limit with β_{odd} and bounds it by $0.3464533951\dots$	The possibility $1/2$ is excluded. Exact β_{odd} and effective convergence remain open.
Large-degree comparison when small prime torsion disappears	Theorem III.99.2 gives the exact joint limit $1/3$.	This is a joint limit, with no quantitative thresholds for d or the least prime divisor.
Source 33, Question 2: optimal cubic coefficient for $s \geq 6$	Corollary II.335.3 determines every inverse coefficient in succession.	The full finite-error envelope at distances other than $1/(2n)$ is not identified.
Source 33, Question 3: equality in the sharp cubic	Theorem II.334.3 classifies all interior equality cases, for every even $s \geq 4$.	The quantitative near-equality part of that question remains open.
Source 33, Question 8: two-direction repair	Theorems II.337.2–II.337.7 give a sharp intersection modulus and optimal repair coefficients through degree $s - 2$.	The theorem concerns full products and an explicit small-error range.
Source 33, Question 1: common-prime quadratic coefficient	The new bracket is $(s - 1)/s^2 \leq c_s^{\text{prime}} \leq (3s - 2)/(2s^2)$.	Both endpoints improve; they still differ.

Prior target	Result here	Remaining qualification
Sources 04 and 17: combine weighted degeneracy control with Fejér amplification	Corollary II.330.5 preserves the retained count and lowers the proposed threshold from L^m to L^{m-1} .	Optimal threshold exponents and uniform treatment of small torsion remain open.

The two exact limiting statements in the norm section have different scopes. In the unrestricted odd-order problem, the limiting discrete dual Γ has no element of order two but can have odd torsion. In the large-prime-divisor limit Γ is torsion-free, so its compact dual $K = \widehat{\Gamma}$ is connected. The latter extra property is exactly what permits the Riesz–Sobolev input. It must not be silently inserted into the former problem.

Similarly, the arrangement classification genuinely uses two-torsion in the target. It explains why the all-group polynomial can have a strictly better second-order behavior on odd targets. Its equality family also requires a quotient of order two in the vertical domain. Neither requirement is compatible with the common-odd-prime example, which is why a separate cancellation analysis is necessary.

V.224.2 What changes at the Gowers interfaces

In Gowers’s Section 9, the restriction theorem supplies a new sufficient size threshold at the same retained-count scale as the Fejér refinement in source 17. It includes an all-group certificate, but its clean prime-uniform corollary assumes a threshold large enough to eliminate short torsion. The graph-energy interpolation used to pass from quadruples to sixteen-tuples is an inherited ingredient, reproduced in (II.330.20); the new step is the translated survival estimate.

In Section 12, respected arrangements become a testing problem near complete agreement on a full product. Source 33 already identified their model class and supplied the coarse decoder and two sharp leading terms. Theorems II.333.2 and II.334.1 retain the constant Fourier coefficient and its tail together, which is what allows the whole extremal polynomial and its equality cases to emerge. These are stronger full-product repair statements. Gowers’s sparse-domain applications require additional control of which arrangements are available and how conditioned queries overlap.

The real norm comparison influences centered cube estimates by changing the admissible relation between U^2 mass and higher uniformity. Our result identifies its asymptotic limit and characteristic dependence. Since the proof does not give a finite-degree rate, it cannot be used to replace a numerical U^4 coefficient. In particular, the separately staged coefficient 4795/832 in source 45 remains a finite-degree input distinct from this article’s result [243].

The full Section 16 interpolation and numerical closure obligations recorded by the repository remain outside the conclusions proved here [246]. A new estimate at one node of the proof must be propagated through all downstream losses before a new bound for $r_k(N)$, the maximum size of a k -term-progression-free subset of $\{1, \dots, N\}$, can be asserted. Modern global bounds likewise involve a complete parameter analysis; see, for example, [2]. No such global improvement is asserted in this manuscript.

V.224.3 A suggested repository integration order

The archive is a standalone research contribution, so its insertion need not alter previously assigned source numbers. A maintainer can register it as the next available source and retain its theorem labels. The following order keeps dependencies visible:

1. Add the norm limit theorem next to source 29’s moment comparison and asymptotic question. Preserve the distinction between the ratio c_d^{odd} and the coefficient in a 2^d th-power inequality.
2. Add the auxiliary-label theorem and arrangement polynomial next to source 33’s local profile. Its $s = 4$ specialization is inherited; the all-order inequality, classification and higher inverse coefficients are the new conclusions. Retain the common-prime bounds as a separate odd-target refinement.
3. Add the two-direction section as a resolution of the local version of source 33’s Question 8. The sum-normalized failure $e = \varepsilon_s^v + \varepsilon_s^h$ and the mixture-normalized failure are not interchangeable without the factor in Corollary II.337.9.
4. Add the translated restriction theorem next to sources 04 and 17. Preserve the explicit counts of repeated tuples and the distinction between the integer principal line and coefficient relations that become dependent only modulo small torsion.
5. Update the research-question ledger using Table V.224.1, preserving the unresolved parts of compound questions rather than marking every associated topic solved.

For formalization, the natural initial targets are the finite auxiliary moment identity, the scalar convex-chord bound, the monotonicity of $E_s(t)/t$, and the weighted rank-two selector estimate. These are finite or elementary analytic statements with explicit interfaces. The compactness theorem needs Haar measure, duality, weak-star compactness and Fourier approximation, so it is a more substantial formalization project. No Lean file has been generated or compiled for these theorems.

[write] **7 October 2026; sources 60–69.** *As integrated here* (7 October 2026): the norm theorem stands beside sources 29 and 63 in Part III (Section III.94), the arrangement and two-direction sections beside source 33 in Part II (Section II.331), the restriction theorem beside sources 04 and 17 (Section II.310), and the questions in Section V.2; nothing is formalized.

V.225 Computational evidence and reproducibility

Source 60, delivered Section 11.

The archive contains the TeX source, the compiled PDF, the vector figure and its plotting source, the verification programs, their recorded outputs, and a source audit. The main article compiles with ordinary pdfLaTeX and standard packages. The proof text is self-contained apart from standard Fourier and compact-group facts, the explicitly stated Christ–Iliopoulou inequality, and background results credited and reproduced where used.

V.225.1 What was checked exactly

For arrangements, the verifier counts accepted balanced relations by integer convolution on $G \times K$. It obtains distance to a model by enumerating the admissible affine slope maps and optimizing a

modal row value. Thus the defect and the model distance are computed by different finite procedures. It checks both the inequalities and the “if and only if” classifications in their stated interior ranges. It covers $s = 4, 6, 8, 16$, repeated queries, binary and nonbinary targets, and cases with nontrivial bilinear maps.

The principal arrangement suite records 16,976 auxiliary-label comparisons and 32,024 arrangement/profile comparisons, including 120 and 12,928 respective equality cases. Separate suites record 494 robust-intersection comparisons, 8,504 two-direction profile comparisons, and 8,656 odd-target comparisons, for 66,654 reported comparisons in total. The coset family is checked for every index $3 \leq n \leq 24$, including even indices. Additional exact evaluations compare the common-prime obstruction formula with generic integer counts at $p = 17$ and with its proved uniform remainder at larger primes.

For restriction, the verifier evaluates the actual survival probability of a tuple by expanding over its distinct selected points and applying exact character orthogonality. This specifically tests the repeated-tuple correction. There are 88 selector cases, including small torsion, and 90 exact coefficient cases checking the signal, the kernel moment, and its alternating-binomial formula. These calculations use integers and rational numbers, not a rounded Fourier transform.

For the norm argument, the finite tests verify the distinct-pair multiplicities, the elementary tail inequality, the factorial simplification, and the explicit cyclic witness. They cannot test an infinite compactness argument. The separate checks of the trigonometric cotangent formula use floating point and are identified as such; the formula’s proof is analytic.

An additional exact polynomial program checks the cancellation of the highest degree, the positive-basis identity for $B_s - F_s$, the gap $C_s - P_s$, and inverse-series composition for every even $4 \leq s \leq 32$. It verifies the exceptional $s = 4$ terms and the stated one-direction and two-direction cubic coefficients with rational arithmetic.

V.225.2 Bounded Fourier extrema on small groups

There is a complete finite optimization behind the norm diagnostic table. The function $f \mapsto \sum_{\chi} |\widehat{f}(\chi)|^4$ is convex on the centered cube

$$\{f \in [-1, 1]^G : \sum_x f(x) = 0\}.$$

For odd $|G|$, every vertex has one zero and equal numbers of $+1$ and -1 . A convex function achieves a maximum at a vertex, and translation invariance permits fixing the zero at the identity. The program enumerates all choices of the positive coordinates and computes

$$\|f\|_{U^2}^4 = |G|^{-3} \sum_h \left(\sum_x f(x) f(x+h) \right)^2$$

with integer arithmetic. There is no heuristic pruning.

The two groups of order nine have different optima, which is a useful diagnostic that the group law, not merely its cardinality, is being used. The cyclic agreement with the explicit witness suggests a separate finite extremal question; it is not extrapolated to all odd cyclic groups in any theorem above.

Group	Vertices after fixing the zero	Exact maximum of $\ f\ _{U^2}^4$
$\mathbb{Z}_3$	2	2/9
$\mathbb{Z}_5$	6	36/125
$\mathbb{Z}_7$	20	106/343
$\mathbb{Z}_9$	70	232/729
$\mathbb{Z}_{11}$	252	430/1331
$\mathbb{Z}_{13}$	924	716/2197
$\mathbb{Z}_{15}$	3432	1106/3375
$\mathbb{Z}_{17}$	12870	1616/4913
$\mathbb{Z}_3 \times \mathbb{Z}_3$	70	56/243

Table V.225.1: Complete bounded, real, centered optimization on the listed groups. The cyclic values coincide with (III.99.29). The table makes no claim about unlisted groups.

V.225.3 How to reproduce the deliverable

The accompanying `README.txt` gives commands for building the PDF and running the checks. The verification programs need only the Python standard library. Regenerating the optional vector figure requires NumPy and Matplotlib; the checked figure is already included, so neither is required to compile the manuscript. Run the checks with ordinary Python, without the optimization flag that disables assertions. All computation described as exact in this section uses integer or rational arithmetic. The mathematical results do not depend on a numerical optimizer or on the finite tests passing.

V.226 Further research questions and a proposed programme

Source 60, delivered Section 12.

The theorems sharpen several questions enough to make the next steps testable. The following programme separates questions that may yield further local improvements from the larger task of transporting those improvements through Gowers’s sparse structural argument.

V.226.1 Fourier extrema and effective compactness

Research question V.226.1 (The unrestricted odd-order limit). Does every centered real function on a finite odd-order abelian group satisfy

$$\|f\|_{U^2}^4 \leq \frac{1}{3} \|f\|_{\infty}^4?$$

By Theorem III.99.1, this is equivalent to $c_{\infty}^{\text{odd}} = 1/3$. The connected compact case is settled here. The remaining issue is the effect of persistent finite odd-order quotients. The mass-cap proof treats two large inverse pairs as independently placeable, even though a bounded function may not realize those coefficient magnitudes simultaneously. A stronger joint coefficient inequality is one concrete route to reducing the upper endpoint.

[write] **7 October 2026; sources 60–69.** *Answered by source 63:* yes, by its Theorem III.96.2, so $c_{\infty}^{\text{odd}} = 1/3$ (item III.1).

Research question V.226.2 (Effective high-degree and high-characteristic thresholds). Find explicit functions $d_0(\varepsilon)$ and, for the connected limit, $P_0(\varepsilon)$ such that

$$c_d^{\text{odd}} \leq B + \varepsilon, \quad C(d, P) \leq \frac{1}{3} + \varepsilon$$

in the respective ranges, where B is the upper endpoint in (III.99.6). The Fourier-tail lemma is already quantitative: a moment order of size ε^{-1} and an initial spectrum of size ε^{-2} make its tail small up to absolute constants. The non-effective step is stabilizing algebraic relations among these frequencies and controlling the smoothing needed to convert finite norm information into a supremum bound. Is there a bounded-height relation model that avoids an unrestricted compactness extraction?

[write] **7 October 2026; sources 60–69.** *Answered by source 63* for both quantities: $d_0(\varepsilon) = d_\varepsilon = 8 \max\{10, \lceil \log_2(10/\varepsilon) \rceil\}$ (its Corollary III.97.6) gives $c_d^{\text{odd}} \leq 1/3 + \varepsilon \leq B + \varepsilon$, and $C(d, P) \leq c_d^{\text{odd}}$ for $P \geq 3$, so no P_0 is needed (checked at the write). The optimal rate is source 63's Research question V.237.2.

Research question V.226.3 (Exact finite cyclic and bounded-exponent extrema). Is the exact bounded Fourier maximum on $\mathbb{Z}_q$, for every odd q , equal to

$$\frac{q^3 - 4q + 3}{3q^3}?$$

The explicit witness attains this value, and exhaustive enumeration confirms optimality for odd $3 \leq q \leq 17$. A proof must control the energy of every centered cube vertex, including at composite orders. More broadly, determine $\beta(q)$ for groups of exponent dividing a fixed odd q , and classify extremizers. At $q = 3$, the finite computation on $\mathbb{Z}_3^2$ and Corollary III.99.10 give the concrete starting interval $56/243 \leq \beta(3) \leq 5/18$.

[write] **7 October 2026; sources 60–69.** *The first question is answered by source 63:* on every abelian group of odd order q the maximum is at most $1/3 - 4/(3q^2) + 1/q^3 = (q^3 - 4q + 3)/(3q^3)$, attained on $\mathbb{Z}/q\mathbb{Z}$ (its Theorem III.96.4); the values agree with the table above for every odd $3 \leq q \leq 17$ (checked at the write in exact rationals). The bound is not attained on $\mathbb{Z}_3 \times \mathbb{Z}_3$ ($56/243 < 232/729$), so the questions on fixed exponent, $\beta(3)$ and the extremizers stay open (source 63's Research question V.237.4; item III.1).

Research question V.226.4 (Stability at the connected limiting constant). For uniformly bounded centered functions with ratio approaching $1/3$ in the joint regime of Theorem III.99.2, can one identify half-circle Bohr sets or a comparably explicit family after passage to an appropriate quotient? The equality and stability theory in [242] supplies structure on a connected compact limit. Transferring that structure back to finite groups, with a quantitative error and an explicit metric, is an additional problem. For initially unbounded functions, the natural first target may be Fourier fourth-power approximation after positive smoothing rather than pointwise or L^1 approximation.

V.226.2 Sharp repair beyond the classified equality cases

Research question V.226.5 (Quantitative near-equality for arrangements). Fix s and keep δ in a compact subinterval of $(0, 1/4)$. If

$$\varepsilon_s(f) - B_s(\delta) \leq \zeta,$$

how close must f , modulo an exact model, be to a classified map $a\mathbf{1}_C(x)\ell_2(y)$? Determine the best exponent of ζ and its dependence on the distance of δ from the endpoints. The proof identifies four

separate losses: collision concentration, convex-chord slack, Fourier-cap slack, and synchronization of active increments. A quantitative argument must make these losses compatible with one coset and one homomorphism. It should also explain how closeness to equality constrains δ to be near $1/(2n)$.

Research question V.226.6 (The common-prime second coefficient). Close the interval

$$\frac{s-1}{s^2} \leq c_s^{\text{prime}} \leq \frac{3s-2}{2s^2}.$$

The lower construction spreads half-interval errors over $\lfloor \sqrt{p} \rfloor$ rows; the upper bound uses the maximal cancellation probability of two errors in an odd-order target. A matching theorem would need to couple the derivative value distribution with the row collision profile. The separate worst cases for those two quantities need not be realizable by the same map. Is the lower endpoint the correct value, or is there another family with more favorable cancellation?

Research question V.226.7 (The finite-error arrangement envelope and activation). Determine the feasible pairs (δ, ε_s) beyond the neighborhood of zero and identify the largest error threshold admitting a canonical repair branch. Even within $0 < \delta < 1/4$, equality in B_s occurs only at $1/(2n)$. Can the universal bound be improved between those distances, perhaps with a piecewise expression sensitive to allowable subgroup indices? The coarse decoder is sufficient to activate our sharp local bound but is not designed to locate a genuine testing threshold. This problem requires structural information beyond formal inversion of a polynomial that is not globally increasing.

Research question V.226.8 (The optimal two-direction profile). Is $P_s(\delta)$ the beginning of an explicit sharp finite-error profile? The sparse-row example has defect $C_s(\delta)$, and (II.337.13) shows that their first possible difference is at order δ^{s-1} . Determine the optimal inverse coefficient at degree $s-1$, beginning with the unresolved cubic for $s=4$. Also classify equality or near-equality in the robust-intersection theorem beyond the displayed binary strip unions. The bilinear-part separation and the distributional intersection argument may have different sharpness mechanisms away from zero.

Research question V.226.9 (Several coordinate directions and higher models). For a product of three or more groups, identify the exact class invisible to a specified family of directional arrangement tests, then prove sharp repair to the compatible multilinear class. The two-direction variance argument suggests using the orthogonal decomposition by coordinate subsets, while the bilinear separation argument suggests a hierarchy of discrete derivative gaps. Can one avoid an exponential accumulation of coarse synchronization constants? An explicit first case is a triaffine model on $G_1 \times G_2 \times G_3$ tested in all three directions.

V.226.3 Restriction, sparse domains, and the global proof

Research question V.226.10 (The best translated positive kernel). For a nonnegative trigonometric polynomial $0 \leq K \leq 1$ of bounded Fourier support, optimize the signal on the principal line relative to its zero coefficient, the total zero-sum coefficient weight, and the repeated-tuple survival correction. Fejér gives the explicit scale L^{m-1} , but the present proof does not establish its optimality. A useful formulation keeps all three quantities in the signed objective $\eta \mathbb{E}X - \mathbb{E}Y$ rather than maximizing signal amplification alone. Can a different kernel improve the retention constant without restoring a larger threshold?

Research question V.226.11 (Uniform restriction with persistent short torsion). Can a finite family of translated selectors eliminate the principal-line leakage measured by $S_{m,L}[q]$ while preserving polynomial retention uniformly over all finite abelian groups? For fixed q , a single Fejér selector has a leakage-to-signal ratio tending to $1/q$, so increasing the bandwidth alone does not solve the problem. A promising route is to separate the low-order quotient, apply a rank-sensitive or exact restriction there, and combine it with a large-order selector. The cost must be controlled in the number of relevant torsion layers rather than in the total group size.

Research question V.226.12 (Sharp domain-size thresholds). Is the exponent $m - 1$ in Corollary II.330.5 necessary for retaining its stated coefficient, or can it be lowered uniformly as the prime modulus varies? The nonprincipal weight and repeated-point correction are bounded independently in the proof. Their worst cases may be incompatible with a graph having many respected relations. An energy-sensitive estimate for the actual repeated bad tuples, or for the weighted rank-two fibers, could improve the threshold further. A negative answer should exhibit a family that obstructs the prescribed retained count, not merely a family on which one particular selector performs poorly.

Research question V.226.13 (From full products to sparse arrangement domains). Which quantitative mixing assumptions on a subset $A \subseteq G \times H$ suffice for a polynomial repair theorem in the conditional arrangement failure rate? The full-product proof uses independent auxiliary labels, a cocycle comparison at every increment, and rowwise integration. A sparse theorem must replace each of these with a controlled witness-count statement. It should record domain density, conditional marginal distortion, and overlap expansion separately. Proving such a theorem is the most direct route from the new sharp local profiles to the sparse structural steps in Gowers's Sections 12–15.

Research question V.226.14 (A complete quantitative propagation and formal interface). After a sparse repair theorem is available, propagate the improved restriction and repair constants through the full interpolation and density-increment chain. Every change should specify the old input, the replacement, the downstream recurrence, and the parameter range in which it improves the final statement. In parallel, formalize the finite auxiliary moment theorem and selector certificate before tackling compactness. A complete proof of a better global bound requires these missing interfaces and numerical closure; the local sharpness results alone do not imply it.

Suggested order of work. Questions V.226.6, V.226.8, and V.226.10 have explicit finite-dimensional objectives and direct diagnostics. Questions V.226.1 and V.226.3 isolate a substantive Fourier extremal problem, with exact examples and a narrow universal interval. Question V.226.2 asks for the quantitative content hidden by the compactness step. Questions V.226.13 and V.226.14 are the larger structural programme needed to influence the global Szemerédi bound. This ordering allows further local discoveries to be validated before they are assigned downstream consequences.

V.227 The inherited coarse decoder, with proof

Source 60, delivered Appendix A.

This appendix supplies the bootstrap used in Theorem II.335.2. The quantitative decoder is the one already proved in the full-product report [248]; it is reproduced so that the sharp polynomial

and its inversion have a self-contained dependency chain. The majority-correction method is in the tradition of Blum, Luby, and Rubinfeld [49]. We retain independent auxiliary labels at every query, including when the group coordinates repeat.

Lemma V.227.1 (Homomorphism correction with labels). *Let U, W be finite abelian groups, let S be a finite nonempty set, and let $F : U \times S \rightarrow W$. Suppose*

$$e = \mathbb{P}_{u,v,\alpha,\beta,\gamma}[F(u, \alpha) + F(v, \beta) \neq F(u + v, \gamma)] < \frac{1}{6}.$$

Then there is a homomorphism $L : U \rightarrow W$ such that

$$\mathbb{P}_{u,\alpha}[F(u, \alpha) \neq L(u)] \leq \frac{e}{1 - 2e}. \quad (\text{V.227.1})$$

Proof. For fixed $z \in U$, consider

$$A_z(y; \alpha, \beta) = F(z + y, \alpha) - F(y, \beta).$$

Its two independent copies agree with probability at least $1 - 2e$. Indeed, introduce one further independent label γ and compare

$$\begin{aligned} F(z + y, \alpha) + F(y', \beta') &= F(z + y + y', \gamma), \\ F(y, \beta) + F(z + y', \alpha') &= F(z + y + y', \gamma). \end{aligned}$$

Each equality fails with probability e : with z fixed, its two group inputs remain independent and uniform. If both hold, the two copies of A_z agree.

Writing $p_z(w) = \mathbb{P}[A_z = w]$, it follows that

$$\max_w p_z(w) \geq \sum_w p_z(w)^2 \geq 1 - 2e.$$

Let $L(z)$ be the unique mode; its mass exceeds $1/2$. Thus

$$\mathbb{P}[A_z = L(z)] \geq 1 - 2e \quad (z \in U). \quad (\text{V.227.2})$$

For fixed z, z' , the identity

$$A_{z+z'}(y; \alpha, \beta) = A_z(y + z'; \alpha, \gamma) + A_{z'}(y; \gamma, \beta)$$

holds pointwise. Each of its three terms equals its mode except on an event of probability at most $2e$. Since $6e < 1$, some choice of the arguments realizes all three modes simultaneously. Therefore $L(z + z') = L(z) + L(z')$.

Finally the test error can also be written as

$$e = \mathbb{P}_{u,v,\alpha,\beta,\gamma}[F(u, \alpha) \neq A_u(v; \gamma, \beta)].$$

Conditional on u, α , if $F(u, \alpha) \neq L(u)$ then the event $A_u(v; \gamma, \beta) = L(u)$ forces a failure. Its probability is at least $1 - 2e$ by (V.227.2), and its labels are independent of α . Hence

$$e \geq (1 - 2e)\mathbb{P}[F(u, \alpha) \neq L(u)],$$

proving (V.227.1). □

Lemma V.227.2 (Affine correction with labels). *For $F : G \times S \rightarrow K$, let $e = e_4(F)$ be the auxiliary test error of Section II.333. There is an affine map $A : G \rightarrow K$ such that*

$$\mathbb{P}_{x,u}[F(x,u) \neq A(x)] \leq 6e. \quad (\text{V.227.3})$$

If $e < 1/6$, the upper bound can be replaced by $e/(1 - 2e)$.

Proof. If $e \geq 1/6$, any affine map satisfies (V.227.3). Otherwise choose an anchor $(x_4, u_4) = (z, w)$ whose conditional test error e_0 is at most e . Conditional on that anchor, set

$$x_1 = z + u, \quad x_2 = z + v, \quad x_3 = z + u + v,$$

with u, v independently uniform. The translated function $Q(u, \alpha) = F(z + u, \alpha) - F(z, w)$ has homomorphism-test error e_0 . Lemma V.227.1 supplies a homomorphism L with disagreement at most $e_0/(1 - 2e_0) \leq e/(1 - 2e)$. Translating back gives the affine map $A(x) = L(x - z) + F(z, w)$. In this range $e/(1 - 2e) < 3e/2 \leq 6e$. $\square$

Lemma V.227.3 (Separation of affine maps). *Two distinct affine maps from a finite abelian group to an abelian group disagree on at least half their inputs.*

Proof. Their difference is $L(x) + c$. If $L = 0$, a nonzero difference never vanishes. If $L \neq 0$, its zero set is empty or a coset of $\ker L$, whose index is at least two. $\square$

Theorem V.227.4 (Coarse full-product repair). *For $f : G \times H \rightarrow K$, if $\varepsilon_4(f) < 1/216$, then*

$$\delta_V(f) \leq 6\varepsilon_4(f) + \frac{36\varepsilon_4(f)}{1 - 72\varepsilon_4(f)} \leq 60\varepsilon_4(f). \quad (\text{V.227.4})$$

Proof. Let e_h be the conditional four-difference failure probability at increment h . Apply Lemma V.227.2 to $F_h(x, y) = D_h f(x, y)$ with label set H . Choose $A_h \in \text{Aff}(G, K)$ such that

$$d_h := \mathbb{P}_{x,y}[D_h f(x, y) \neq A_h(x)] \leq 6e_h.$$

Put $d = \mathbb{E}_h d_h$, so $d \leq 6\varepsilon_4(f) < 1/36$. The derivative identity

$$D_{h+k} f(x, y) = D_h f(x, y + k) + D_k f(x, y)$$

and a union bound yield

$$\mathbb{P}_{h,k,x}[A_{h+k}(x) \neq A_h(x) + A_k(x)] \leq 3d. \quad (\text{V.227.5})$$

For each h, k , the compared affine maps are identical or differ on at least half of G , by Lemma V.227.3. Thus

$$\eta := \mathbb{P}_{h,k}[A_{h+k} \neq A_h + A_k] \leq 6d < 1/6.$$

The affine maps form a finite abelian group under pointwise addition. Apply Lemma V.227.1, with a singleton label set, to the map $h \mapsto A_h$. It produces a homomorphism

$$\mathcal{A} : H \longrightarrow \text{Aff}(G, K)$$

whose disagreement proportion is

$$q \leq \frac{\eta}{1 - 2\eta} \leq \frac{6d}{1 - 12d}. \quad (\text{V.227.6})$$

Set

$$\ell(y) = \mathcal{A}(y)(0), \quad B(x, y) = \mathcal{A}(y)(x) - \mathcal{A}(y)(0).$$

Then ℓ is additive and B is biadditive. For the residual $r(x, y) = f(x, y) - B(x, y) - \ell(y)$,

$$\mathbb{P}_{h,x,y}[D_h r(x, y) \neq 0] \leq d + q.$$

For each x , choose a most frequent value $a(x)$ of $r(x, \cdot)$. Since independent row samples disagree with probability at least one minus the largest atom, averaging the row collision identity gives

$$\mathbb{P}_{x,y}[r(x, y) \neq a(x)] \leq d + q.$$

This constructs a member of $\mathcal{V}$ with the required distance. Using $d \leq 6\varepsilon_4(f)$ in (V.227.6) gives the first bound of (V.227.4). Finally, $1 - 72\varepsilon_4(f) > 2/3$, so the fraction there is at most $54\varepsilon_4(f)$, proving the second bound. $\square$

Corollary V.227.5 (The exact kernel of the test). *For every even $s \geq 4$, $\varepsilon_s(f) = 0$ if and only if $f \in \mathcal{V}$.*

Proof. Model functions always pass. Conversely, the moment identity gives $\varepsilon_4(f) \leq \varepsilon_s(f) = 0$, and Theorem V.227.4 gives distance zero to $\mathcal{V}$. $\square$

V.228 Exact computational checks and constructive content

Source 61, delivered Section 10.

The accompanying Python programs implement the finite constructions and run exact checks. They do not serve as a proof assistant, and exhaustive verification of small cases is not used in place of any general argument above. Their purpose is to catch sign errors, torsion exceptions, boundary mistakes, and normalization errors in the executable versions of the proofs.

V.228.1 Representation and arithmetic

A group is represented as an explicit product $\prod_i \mathbb{Z}/q_i\mathbb{Z}$. Elements are tuples of canonical integer representatives. A character is indexed by a tuple a , with value

$$\chi_a(x) = \sum_i \frac{a_i x_i}{q_i} \pmod{1}.$$

Let $Q = \text{lcm}_i q_i$. The character value is stored as a numerator v modulo Q . The open radius-1/5 test is the integer comparison

$$5 \min(v, Q - v) < Q.$$

The color index is $\lfloor 5mv/Q \rfloor$. No floating-point trigonometry or approximate Fourier transform is used.

Additive energies are computed by exact convolution of counting functions. The BSG score and its thresholds use rational arithmetic. For the weighted tests, each complex weight is a Gaussian integer stored as a pair of integers; convolution and squared magnitudes are computed by integer operations. Only the optional elapsed-time field in the run record is inexact.

V.228.2 Coverage of the delivered run

The run recorded in `checks/results.json` completed with every assertion passing. The principal coverage is:

Check family	Cases
Root-grid inequality for every order $2 \leq q \leq 10000$	9,999
Nonzero-element character distributions in nine small groups	40
All forbidden sets in eight groups, at orders 1, 2, 3, 8	1,304
All full-domain maps in five domain/codomain pairs	1,826
Those maps tested at orders 2, 3, 8	5,478
BSG cores for all maps $\mathbb{Z}/4\mathbb{Z} \rightarrow \mathbb{Z}/4\mathbb{Z}$	256
Largest-cell size and energy checks for those cores	256
Sampled partial maps in mixed-torsion groups	100
Gaussian-integer weighted exactness and retention checks	200

The five exhaustive map families are

$$\mathbb{Z}/2 \rightarrow \mathbb{Z}/2, \quad \mathbb{Z}/3 \rightarrow \mathbb{Z}/3, \quad \mathbb{Z}/4 \rightarrow \mathbb{Z}/4, \quad \mathbb{Z}/5 \rightarrow \mathbb{Z}/3, \quad \mathbb{Z}/4 \rightarrow \mathbb{Z}/6.$$

Their respective numbers of maps are 4, 27, 256, 243, 1296. The avoidance tests use every forbidden subset, not only symmetric subsets, in

$$\mathbb{Z}/2, \mathbb{Z}/3, \mathbb{Z}/4, \mathbb{Z}/5, \mathbb{Z}/6, \mathbb{Z}/8, (\mathbb{Z}/2)^2, \mathbb{Z}/2 \times \mathbb{Z}/4.$$

The sampled partial-map family has domain $\mathbb{Z}/4 \times \mathbb{Z}/2$ and codomain $\mathbb{Z}/3 \times \mathbb{Z}/4$. Sampling uses seed 20261006 and is supplementary to the exhaustive cases.

Five torsion-obstruction instances were also checked, including every two-point subset in the indicated coloring groups:

e	r	$m = 2^{e-1}$	$ H $	Required colors
2	1	2	4	4
2	2	2	16	16
3	1	4	8	8
3	2	4	64	64
4	1	8	16	16

The canonical-section graph difference and energy formulas were verified for the same parameters. Every two-point domain restriction failed the required Freiman condition, as predicted. The inequality $40^{17} < 2^{91}$ and every arithmetic identity in the order-eight constant calculation were checked with integers.

V.228.3 How to reproduce the checks

From the package directory, run

```
python3 code/run_checks.py
```

The program requires Python 3.10 or later and uses only the standard library. It rewrites `checks/results.json` and prints the same record. The test suite must be run with assertions enabled; do not use Python's `-O` option.

The main construction can also be called directly:

```
from exact_restriction import AbelianGroup, freiman_partition
G = AbelianGroup((8,))
H = AbelianGroup((16,))
phi = {x: x for x in G.elements()}
characters, cells, defects = freiman_partition(G, H, phi, 8)
```

This example returns singleton cells, illustrating the order-eight obstruction. The module must be on the Python import path, for example by running from the `code/` directory.

V.228.4 Constructive versus efficient

The proofs give finite algorithms: enumerate characters, repeatedly minimize the number of remaining unseparated defects, color by exact interval indices, and select a largest cell. The BSG proof similarly maximizes a finite score and prunes a finite neighborhood. However, the supplied implementation explicitly enumerates the group, graph sumsets, and character table. It is intended for small-instance certificates, not as a polynomial-time algorithm in the bit length of a compressed group presentation.

The guarantee on the *number of colors* is independent of the ambient group size. That mathematical guarantee must not be confused with an algorithmic running-time guarantee. Computing the defect set itself may dominate the cost.

V.229 Integration and proof-assistant targets

Source 61, delivered Section 11.

V.229.1 Recommended mathematical placement

A suitable new research directory is

```
Combinatorics/Ramsey/Research/GowersSzemerédi/
torsion-exact-restriction/.
```

The article is self-contained and need not be merged into the existing consolidated manuscript before review. Its labels have the prefix `ter:` to reduce collision risk in a later integration. The package includes source, PDF, exact checking programs, source notes, and a formalization plan. No repository files were changed in preparing this delivery.

[write] 7 October 2026; sources 60–69. The suggested directory was not used: by its placement (39d3295c4) this source is source 61 of this report, with labels `gsr:fm:61:`, `gsr:en:61:` or `gsr:fz:61:` in front of its `ter:` labels, and files under the prefix `61-exact-restriction-`.

The most useful dependency order is

root-grid estimate $\longrightarrow$ defect separation $\longrightarrow$ avoidance partition,
 Plünnecke–Ruzsa $\longrightarrow$ vertical budget $\longrightarrow$ exact graph partition,
 BSG core + exact graph partition + sumset energy $\longrightarrow$ threshold-free order-eight theorem.

The obstruction theorems are logically independent of the BSG estimate and can be integrated as a separate subsection.

V.229.2 Minimal formal interfaces

There are two reasonable formalization routes. One can work from the start with arbitrary finite abelian groups and their character groups, or first implement the construction for explicit products of cyclic groups and then transport it across an isomorphism. The latter route mirrors the integer certificate program and isolates the group-classification dependency.

The first target should be a finite separation lemma, not the final numerical corollary. Its abstract input is a finite collection of tests such that every nonzero defect passes at most half of them. Greedy halving then needs only finite sums and cardinalities. Character orthogonality or the root-grid calculation supplies the input later.

The next targets are the interval-width implication and the graph-translate injection. The interval argument must use half-open cells and the strict radius condition; replacing either endpoint convention without rechecking the proof is unsafe. The graph injection is purely algebraic and should not require Fourier analysis.

For relation counts, establish an explicit equivalence between the ordered tuple formulation in Definition II.322.1 and the repository’s `FreimanHom` interface. The counting conventions allow repetitions and count ordered tuples. Next identify $\mathcal{E}_8(B)$ with `additiveTupleCount 8 B` and the respected energy with `phiAdditiveTupleCount 8 B phi`. Exactness then implies equality of those two counts, and the `GammaHomOfOrder` conclusion follows by a scalar inequality.

Finally, the implication to the inspected `lemma_9_3` should choose $N_0 = 1$, derive $\alpha \leq 1$ from the energy cap, and use the integer exponent comparison in Corollary II.328.1. This is a proposed proof plan, not a checked Lean implementation. No sorry-bearing file is shipped as a purported proof, and no formalization-status ledger should be updated solely on the basis of this article.

V.229.3 What requires review before promotion

The main mathematical review points are the universal character estimate in all small torsion orders, the strict interval-boundary argument, the disjoint vertical translates, the distinction between the original graph size and the extracted-core size in the BSG ratio, and the use of ordered repeated-variable relations. The algebraic derivations of the constants are supplied separately from the conceptual claims to make review modular.

A later numerical refinement should use Proposition II.327.2 rather than silently replacing the conservative BSG constants. A later global application must additionally prove the analytic compatibility hypotheses that it uses. Neither numerical improvement nor compatibility follows merely from the existence of an exact restriction.

V.230 Further research questions

Source 61, delivered Section 12.

The questions in this section arise from the present results. They are proposed research problems, not assertions that the same formulations are established open problems in the literature. Their priority and their relation to other versions of additive coloring should be investigated independently.

Question V.230.1 (Determine the exact avoidance exponent). *Is*

$$b_m = \lfloor \log_2 m \rfloor + 1 \quad (m \geq 2)?$$

If not, determine the least m for which the 2-primary examples do not give the optimal universal exponent.

Theorem F.66.3 confines the uncertainty to an additive constant. A proof of the displayed formula would require a mechanism better than the present independent narrow-arc construction, or a more efficient way of using torsion in the separation step. Conversely, a lower bound larger than the displayed value would require a genuinely different obstruction. Exact searches for small (H, D) can be organized as hypergraph coloring problems, with hyperedges consisting of supports of forbidden balanced relations.

Question V.230.2 (The graph-extraction exponent). *Let a_m^* be the infimum of exponents a for which a constant $c_m > 0$ guarantees an exact order- m restriction of size $c_m K^{-a} |B|$ from every graph of difference ratio at most K . What is the asymptotic growth of a_m^* ?*

The present results give a lower bound of order $\log m$ and an upper bound of order $m \log m$. Is $a_m^* = O(\log m)$, or is there a graph obstruction forcing a larger scale? Resolving this requires more than optimizing the constant 5 in the character partition.

Question V.230.3 (Use the structure of vertical defects). *Can the worst-case estimate $|\mathcal{D}_m(\Gamma)| \leq K^{2m+1}$ be replaced by a substantially smaller effective parameter for graph restriction, without assuming that the codomain has prime exponent?*

The defect set is not arbitrary: it is a vertical fiber of a structured iterated difference set. Its additive structure, annihilators, or distribution across primary components may permit a smaller family of separating tests than its cardinality suggests. General-group structural results such as Green–Ruzsa [249] provide context, but no quantitative deduction of this kind is proved here.

Question V.230.4 (Optimal relation-energy powers). *For fixed m , what are the best powers p and constants c_m in a universal exact restriction theorem of the form*

$$\mathcal{E}_m^\phi(B) \geq \alpha |B|^{2m-1} \implies \mathcal{E}_m(B') \geq c_m \alpha^p |B|^{2m-1}?$$

Can the current $O(m^2 \log m)$ exponent be reduced in order of magnitude?

There are at least three separate losses to analyze: conversion to ordinary additive energy, selection of a small-difference graph core, and concentration of a high moment in one exact cell.

Proposition II.327.2 separates the second from the third. A weighted BSG theorem preserving an appropriate relation distribution might avoid the loss incurred by selecting solely through cardinality.

Question V.230.5 (Exact versus approximate tradeoffs). *For arbitrary composite cyclic groups, how much better can relation-energy retention be if exact respect is relaxed to a proportion $1 - \eta$? Can one obtain a useful bound that interpolates between prime-field approximate restriction and the torsion-uniform exact partition?*

The exact theorem removes η entirely, but pays a large power of α . An approximate theorem need not inherit the singleton obstruction in the same form. Identifying the correct torsion-sensitive statistic could give sharper cyclic bounds than a general-group theorem.

Question V.230.6 (Retain the analytic data needed later in Gowers's proof). *Can one construct exact cells while retaining a controlled part of a family of derivative Fourier coefficients and the compatibility of the associated cross-sections, rather than a single prescribed scalar correlation?*

A useful positive result must state the family norm being retained, the common restriction on which it survives, and the consistency conditions across indices. Proposition II.325.4 alone is insufficient. This is the most direct route by which the present local argument might become relevant to later high-degree steps, but no such bridge is proved here.

Question V.230.7 (Cyclic groups versus arbitrary finite abelian groups). *Does restricting both domain and codomain to cyclic groups permit smaller universal exponents than a_m^* or the corresponding energy exponent for arbitrary groups?*

The strongest lower-bound families in Section II.329 increase the rank while fixing the exponent. They therefore do not automatically settle the cyclic optimization. A cyclic analysis must account for primary factors and carries without importing a rank-based obstruction that its hypotheses exclude.

Question V.230.8 (Algorithms with compressed input). *Can a short separation certificate and a large exact cell be found in time polynomial in a natural compressed description of B , ϕ , and the ambient groups, assuming the requisite growth or energy promise?*

The first obstacle is obtaining or representing the relevant defect set without enumerating m -fold sumsets. One possible goal is a certificate verifier whose complexity depends on the certificate and the input rather than on $|H|$. The supplied program establishes finite constructivity only.

Question V.230.9 (Stability of the extremal torsion examples). *If $\chi_m(H, D)$ is close to the largest power permitted by a conjecturally optimal exponent, must (H, D) contain a large quotient or subsystem resembling the p -primary examples? Is there a corresponding stability statement for nearly singleton-only graph restrictions?*

Such a theorem would turn the present examples from lower-bound witnesses into a structural explanation of difficult instances. It would also indicate which extra hypotheses could genuinely improve the restriction theorem, as opposed to merely improving its constants on easy cases.

V.231 Conclusion

Source 61, delivered Section 13.

The exact-partition method yields an elementary, torsion-uniform route from graph growth to bounded-order additive rigidity. Its central quantitative feature is the separation of a finite

vertical-defect budget by $O(\log |D|)$ characters. This supplies a complete partition, preserves weighted energies with explicit losses, and turns a conservative energy-to-growth input into a threshold-free exact strengthening of the stated Lemma 9.3 conclusion.

The coloring obstruction determines the leading logarithmic dependence of the universal defect-coloring exponent on relation order. The graph obstruction establishes a distinct lower bound and shows why one should not extrapolate prime-modulus intuition to arbitrary torsion. The substantial remaining gap in graph exponents and the missing analytic-compatibility bridge are explicit. These are the natural next mathematical targets; an improved global Szemerédi bound is not a consequence of the present manuscript.

V.232 Source and validation ledger

Source 61, delivered Appendix A.

The original Gowers paper was inspected at its Lemma 7.5 and Lemma 9.3 statements and their surrounding arguments. The corresponding original PDF pages include printed pages 505–506 and 516. The repository inspection used the pinned commit recorded in Section F.66; the source notes in the package record the relevant file identifiers.

The principal inspected repository resources were `Sections06_07.lean`, `Sections08_09.lean`, `Proofs09Restriction.lean`, and the README of the local quantitative refinements collection. The last resource is a research guide, not a proof certificate. Its record of stronger prime-field constants is used to delimit the present contribution, not as a premise in any proof.

The mathematical validation consists of the written proofs, exact finite computations, and consistency checks of the constants. It does not include human referee review, an independent proof-assistant run, or a full compile of the PROVEIt Lean development. The Python files are reference implementations; the tests have finite coverage, explicitly described in Section V.228. The LaTeX source and its bibliography are original to this delivery; no copy of a cited paper is included.

V.233 Exact checks and reproducibility

Source 62, delivered Section 10.

The package includes `code/verify.py`, which uses only the Python standard library. Fractions and integer arithmetic are used for all asserted identities and inequalities. Decimal table entries are displays of exact quantities, except that logarithms in the parameter table are rounded floating-point evaluations.

The recorded run checks the following families.

1. **Primitive profiles.** For $m \in \{4, 6, 8, 16\}$, $2 \leq L \leq 32$, and $2 \leq q \leq 160$, all 19 716 response values are checked for nonnegativity and the sampling bound $B_{m,L}(q) \leq q/\varphi(q)$. The dyadic elementary envelope and the large- q bound are also checked wherever applicable.
2. **Weighted arithmetic.** The exact divisor formula and K_m bound are tested in 2 184 modulus cases, including composite and prime moduli. For $L = 2, 3$ and $m = 4$, direct enumeration of coefficient vectors independently checks the divisor formula in 22 cases. The expressions (II.317.9)–(II.317.10) and the finite-length overshoot at $m = 16$ are checked exactly.

3. **Random subset selection.** Six small examples with distinct or equal composite target and domain moduli are evaluated exactly. Joint inclusion probabilities are computed by finite Fourier convolution. Möbius inversion then obtains the probability of every selected subset, and verifies positivity, total mass one, the weighted main-term estimates, and the repeated-vertex correction. These tests use $N \leq 7$ and include all 312 subset outcomes across the six examples.
4. **Parameters and obstructions.** The dyadic recipe, the coarse polynomial inequalities, and the comparison with (II.319.6) are checked on 28 parameter cases, including the three table rows. Primorial examples check the primitive lower bound supplied by the two residues ± 1 . Separate order-two calculations display the all-scalar saturation and the growing unit-dilation gain.

These are finite consistency checks, not proofs of universally quantified statements. The input size threshold is not claimed to be necessary, and the small examples are not advertised as numerical instances saturating it. The universal assertions rest on the preceding arguments.

Reproduction. From the package directory, run

```
python3 code/verify.py
latexmk -pdf -interaction=nonstopmode -halt-on-error article.tex
```

The verifier writes `data/verification.json`. A Makefile provides `make check`, `make pdf`, and `make clean`.

V.234 Repository integration and proof obligations

Source 62, delivered Section 11.

The inspected repository reference is

17123532a2de638477dce1c92eb0b0b20a247d2f.

The concrete interface is `Combinatorics/Ramsey/Lean/GowersSzemerédi/Sections08_09.lean`, whose inspected blob is `0fb6e0a1de3f046f436747603f7eeb06f8d1f074`. The source audit records the additional inspected paths and the limits of the literature comparison.

Suggested mathematical declarations. A formal development could separate the proof into the following interfaces:

Primitive detector.

Nonnegativity of $P_{m,L}$, the rational Ramanujan formula, the grid-sampling inequality, and the elementary envelope H_{el} .

Weighted finite-ring count.

The kernel cardinality (II.317.7), the exact divisor formula (II.317.2), and the rational upper bound K_m .

Random restriction.

A finite probability space of parameters and pointwise membership indicators, with survival products indexed by the distinct support of each tuple.

Objective extraction.

The deterministic implication $\eta G' - (1 - \eta)Y' > 0 \Rightarrow G' \geq (1 - \eta)(G' + Y')$, and the finite averaging principle.

Old-interface corollary.

Corollary II.319.2, followed by the elementary comparison with (II.319.6). The Rosser–Schoenfeld refinement is unnecessary for this step.

No untested Lean file is included. The mathematical statement is deliberately separated from any future implementation decisions about Fourier transform normalization, finite probability spaces, or totient lemmas. The repository’s existing proof and its status ledger are left unchanged.

What this does not propagate automatically. The original-domain Fourier spectrum is not asserted to be preserved by the restriction. The weighted gcd computation concerns one balanced additive relation, not the whole family of arrangement equations in later sections. The proof therefore does not remove the need for the subsequent Bohr-set, multifunction, or higher-order phase arguments. Obtaining a global Szemerédi improvement would require explicit control through those interfaces, not merely substituting one better local exponent.

V.235 Further research questions

Source 62, delivered Section 12.

The following questions are specific continuations of the proved results. They are not claims that the corresponding problems are open throughout the literature; they identify the unresolved tasks in this development.

Question V.235.1 (Can a different detector eliminate the double logarithm?). *Find a bounded nonnegative restriction mechanism that, for every cyclic modulus, gives retention $c_m \alpha^m \eta^{m-1} |B|^{m-1}$ with no logarithmic denominator and a polynomial threshold. Theorem II.320.2 shows that a proof cannot obtain this merely by replacing the uniform bound for the current primitive response by an absolute constant.*

Question V.235.2 (The sharp weighted degeneracy constant). *Determine $\sup_{L \geq 2} \kappa_{m,L}^*$. The limit is known from Theorem II.317.1, but at $m = 16$ it is not the answer. Computations for $2 \leq L \leq 40$ are consistent with a maximum at $L = 3$ for $m = 16$; no such global maximum is proved here. A proof should exploit the exact residue distribution in (II.317.6), not substitute the limit for a uniform bound.*

Question V.235.3 (Matrix-valued weighted degeneracy). *Replace the single equation $\sigma \cdot x = 0$ by the matrix of constraints for a higher-dimensional Gowers arrangement. Develop a weighted Smith normal form analogue of (II.317.8), retaining coefficient weights before estimating torsion indices. This is the most direct route from the present method toward the later arrangement-selection lemmas.*

Question V.235.4 (Preserving the original Fourier data). *Can the restriction be chosen to satisfy a useful compatibility condition with specified large Fourier coefficients of the original domain? One target is a common frequency set for a family of restrictions, with polynomial loss in both the number of prescribed coefficients and their threshold. The current proof supplies no such compatibility theorem.*

Question V.235.5 (Noncyclic target groups). *For a general finite abelian target, replace primitive-unit averaging by a family of characters that separates every nonzero defect without losing a factor depending badly on the target rank. The difficulty is visible already for an elementary abelian target: no single scalar-unit orbit reproduces the primitive cyclic calculation.*

Question V.235.6 (Optimal density and accuracy exponents). *Determine whether the powers $\alpha^m \eta^{m-1}$ in the refined retention bound are optimal under the sole hypothesis $S_m(B, \phi) \geq \alpha |B|^{m-1}$. Separate an obstruction to the theorem itself from an obstruction to one-stage Bernoulli filtering. In particular, the primordial detector example does not answer this question.*

Question V.235.7 (Useful derandomization). *Develop an implementation of the conditional-expectation procedure whose running time is controlled by the sparse input size rather than an exhaustive list of all additive tuples and all unit dilations. A first concrete target is fixed $m = 4$ with a compressed difference table and exact arithmetic.*

Question V.235.8 (Formalization and downstream parameter accounting). *Formalize the elementary route first, then propagate its explicit constants through a clearly identified downstream chain. The required output is a new final theorem with its hypotheses and threshold, not a collection of locally improved estimates. The exact weighted gcd identity is an especially small and reusable first formalization target.*

V.236 Conclusion

Source 62, delivered Section 13.

The composite-modulus difficulty in the Fejér restriction argument has two separate causes: low-order target defects and additional short domain relations. Unit dilations address the first, while a weighted gcd count addresses the second. Together they yield polynomial losses for arbitrary cyclic moduli, including an explicit strengthening of the order-eight restriction interface.

The analysis also produces exact arithmetic information that is not visible in an unweighted exceptional-set bound: the maximal weighted error at each filter length, its extremizing divisibility condition on the modulus, its zeta-ratio limit, and a finite-length overshoot. These results delimit both the strength and the limitations of the method. The next substantial step is to extend the weighted arithmetic calculation to the arrangement systems and spectral compatibility conditions used later in the proof.

V.237 Further research questions and a proposed program

Source 63, delivered Section 7.

The two asymptotic problems now have exact first-order answers. The next questions concern rates, extremizers, the efficiency of the intermediate certificates, and propagation through the surrounding argument. The questions below are deliberately separated from the proved theorems. Each includes a concrete mathematical target and a reason that the present proof does not already answer it.

V.237.1 Extremal uniformity norms

Research question V.237.1 (Exact constants at fixed order). Determine c_d^{odd} for one fixed $d \geq 4$, beginning with $d = 4$. Is its supremum attained on a finite group, or only approached along groups of unbounded order?

The inspected reports give $c_3^{\text{odd}} = 1/\sqrt{2}$ and, using the later fourth-order refinement,

$$0.5587725 \leq c_4^{\text{odd}} \leq (832/4795)^{1/4} = 0.6454070 \dots \quad [258].$$

The lower decimal is a truncated lower estimate from the construction recorded in the source collection. Our uniform high-order proof sacrifices too much to compete at order four. A useful approach is to retain the complete autocorrelation moment expansion, including torsion relations, while adding the pointwise doubling constraint from Lemma III.96.1. Neither positive definiteness alone nor the single bound on the largest Fourier coefficient captures all those restrictions.

Research question V.237.2 (The optimal rate of approach to one third). Determine the asymptotic size of $c_d^{\text{odd}} - 1/3$. Can the upper bound $O(2^{-d/8})$ be replaced by a matching upper and lower asymptotic estimate?

There are two distinct tasks. The finite parameter theorem can be optimized by balancing the tail cutoff, smoothing width, and Bohr set loss. That may improve an upper rate without revealing the true one. Independently, a lower-bound family must compare U^2 with U^d , rather than just with L^∞ . The circle square wave and large odd cyclic approximations are natural test families: their bounded U^2 value is known exactly, while the high-order cube integral retains information discarded by the supremum-norm bound. Estimating that integral as d grows could distinguish an optimal rate from an artifact of smoothing.

Research question V.237.3 (Equality and stability for the bounded inequality). Classify equality and near-equality in $\|f\|_{U^2}^4 \leq \|f\|_\infty^4/3$ on compact abelian groups with surjective doubling, and in the finite-order correction on each finite odd-order group.

The compact circle square wave is an equality case; finite cyclic square waves attain the order-dependent optimum. The defect identity (III.96.8) gives an exact necessary condition on autocorrelations. Turning it into structural information about f requires controlling where its second factor is small and then reconstructing a function from near-extremal correlations. A first target is a quantitative theorem on the circle, with an explicit distance and allowed symmetries. A separate target is to decide which finite abelian groups attain the cyclic order- N value. The finite theorem alone does not identify those groups.

Research question V.237.4 (Dependence on group structure). For a prescribed finite odd-order group G , determine

$$b(G) = \sup_{\substack{f: G \rightarrow \mathbb{R}, \mathbb{E}f=0 \\ f \neq 0}} \frac{\|f\|_{U^2}^4}{\|f\|_\infty^4}.$$

More generally, determine the limiting constants when the supremum is restricted to groups of a fixed exponent or to a fixed family of products.

The fixed-group result in Section III.97 shows that $b(G)$ is exactly the limit of $c_d(G)$. This turns a question about arbitrarily high cube order into a finite bounded optimization problem. Cyclic

groups are settled. Products of small odd cyclic groups provide a tractable next test: they contain many Fourier relations, and the cyclic sharpness construction need not transfer with the same order-dependent value. Exact arithmetic and symmetry reduction can guide conjectures, but a finite search must not be substituted for a bound over all product dimensions.

Research question V.237.5 (Other reality and centering constraints). What are the analogous high-order constants for complex functions with a prescribed phase symmetry, or for real functions with a specified nonzero mean and range?

The assumptions in the present theorem are substantial. A nontrivial complex character is centered and has every U^d norm equal to one, so the unrestricted complex comparison constant is one. An even-order group has a real character giving the same obstruction. Meaningful extensions therefore need an explicit constraint that excludes these examples. For nonzero mean, the doubling argument still supplies information, but its average contains the extra term $2(\mathbb{E}f)^2$ and sharpness may depend on torsion. Determining the right constrained problem should precede optimization of its constants.

V.237.2 Optimal positive covers and phase removal

Research question V.237.6 (The exact positive-cover cost). For $B = [0, \ell]^k$ and $0 < w \leq \ell/2$, minimize $\mu(B)$ over nonnegative measures supported on B subject to

$$\int_B (w - \|a - c\|_1)_+^2 d\mu(c) \geq v_k w^{k+2} \quad (a \in B).$$

Is the face measure of Theorem IV.68.2 optimal in any nontrivial regime?

The convolution identity is exact at corners and has positive excess at many other targets. That excess may permit a smaller or nonconstant measure. A dual certificate would be especially useful: find a positive measure λ on target points for which $\int K_w(a - c) d\lambda(a) \leq 1$ for every center c ; integrating the covering constraint then gives a lower bound for the required mass. Symmetry reduces the number of parameters in small dimensions, but a proposed optimum must satisfy the inequalities on the whole box, including intersections of boundary layers.

Research question V.237.7 (The next term of the optimal local coefficient). Determine whether

$$\lim_{k \rightarrow \infty} k \left(1 - \frac{c_k^{\text{loc}}}{v_k} \right)$$

exists, and compute it if it does. In particular, is the optimal relative deficit of order $1/k$, or is it $o(1/k)$?

The proved certificate gives $1/D_k = 1 - 2/(3k) + 14/(9k^2) + O(k^{-3})$ and therefore an upper bound $2/3$ for a possible limiting deficit. The geometric ceiling only gives its nonnegativity. These two facts do not identify a second-order optimum. Even a better cover may leave slack in the window Cauchy–Schwarz or mixed Gowers inequalities. A matching upper example must be an admissible analytic input to phase removal, not merely a difficult target for one chosen covering construction.

Research question V.237.8 (Exact local constants at small degree). Determine c_k^{loc} at $k = 1$ or $k = 2$ under the normalization of (IV.68.2).

The new and previous covers agree at $k = 1$, where the retained fraction of the ceiling is $9/11$. They do not prove that this is the optimal local coefficient. In small degree one can study the phase and interval geometry directly, potentially bypassing one of the general-purpose averaging steps. This would also test whether optimizing positive covers is the correct route to the analytic optimum or only to a better sufficient estimate.

Research question V.237.9 (Lattice corrections and simultaneous limits). Improve the exact finite coefficient $R_{k,m,N}$, and determine whether a different localization construction can attain the geometric scale under a weaker relation between m and k .

The joint-limit corollary in Section IV.68 already identifies the transition for the present certificate; in particular, $m/k^2 \rightarrow \infty$ suffices for $R_{k,m,N}/v_k \rightarrow 1$. Its proof balances the factors $m^k/(m-1)^k$ and $(w/(L+1))^{k+2}$ quantitatively. A better construction could alter that transition. Direct lattice kernel sums, exact interval cube polynomials, and adjusted endpoint weights may improve the finite estimate. This is a concrete route toward stronger finite estimates suitable for use in an iterated argument. The geometric upper estimate also needs uniform control when the degree varies.

Research question V.237.10 (Other windows, boxes, and kernels). Extend the face identity to unequal side lengths, anisotropic ℓ^1 kernels, or other localization shapes, and determine which extensions lead to an improved analytic localization theorem.

The positive-power theorem already separates the geometric identity from its square-kernel application. Unequal coordinate scales are natural because a multilinear box need not have equal widths in other settings. A geometric generalization is useful to the analytic problem only when an appropriate window-count identity produces its kernel. Thus both pieces should be proved together: a positive cover with controlled mass, and an exact or lower-bounded window multiplicity for the relevant configurations.

V.237.3 Counting applications and proof infrastructure

Research question V.237.11 (A joint density–uniformity expansion). Determine the optimal cube-count upper envelope when the density δ , cube dimension d , and uniformity size $u = \|f\|_{U^d}$ vary together.

The quartic coefficient is now asymptotically sharp as a function of d , but the remainder has combinatorial factors involving 2^d . Fixed- d small-perturbation asymptotics alone do not describe a regime in which d also grows. A useful next theorem would specify an explicit relation between u/δ and d under which the quartic term dominates the remainder, and then exploit the actual fifth and higher coefficients instead of bounding all of them by mixed Gowers–Cauchy–Schwarz.

Characteristic-dependent dependencies among cube vertices become relevant at those higher orders.

Research question V.237.12 (Propagation through the density-increment argument). Track the new local coefficient through the later induction in Gowers’ paper with explicit constants and every scale condition. Which quantitative bottleneck becomes dominant after this replacement?

The improvement from carry-pattern losses to factorial-size local coefficients is substantial at that step. A final theorem about progression-free sets also depends on the extraction of a multilinear piece, transfer between progressions and cyclic groups, the polynomial approximation step, and the length lost under iteration. An honest propagation calculation must preserve those dependencies. A useful deliverable would be a symbolic recurrence with independently verified hypotheses at each transition, even if the final asymptotic shape remains unchanged.

Research question V.237.13 (Formal verification of the new interfaces). Formalize the bounded inequality, the exact face-cover identity, and the finite parameter comparison theorem in the ProveIt environment, preserving their stated normalizations and quantifiers.

The bounded finite-group theorem is a suitable first target: it uses a scalar inequality, translation invariance, bijectivity of doubling, and a finite sum identity. The exact finite-order correction adds a finite-dimensional extremal argument and a direct cyclic calculation. The norm-limit theorem has a longer dependency chain involving Fourier support, probability kernels, and Bohr sets; the positive-cover interface additionally requires polynomial polarization and mixed cube inequalities. Appendix V.238 suggests an order of integration so that the mathematical interfaces can be checked before the asymptotic assembly.

Among these questions, the bounded equality problem, the optimal cover in dimensions one and two, and explicit two-parameter localization bounds appear to be the most direct next projects. The exact fourth-order constant and the optimal high-order rate require additional extremal ideas beyond the current proofs.

V.238 Source audit, verification, and integration

Source 63, delivered Appendix A.

V.238.1 Inspected sources and comparison baseline

The primary paper is Gowers’ 2001 article [1]. We inspected its original PDF and the corrected transcription in the ProveIt Ramsey directory. The repository was retrieved at commit

`d7b83fa1fe60478e3ed8525d58f747a1d76e35c2`.

The inspected research index lists Sources 01–46. Its combined TeX manuscript contains the first 33 reports; later source reviews and integration notes were inspected separately where they affect the present comparison. These are different layers of evidence: an entry in the index is not by itself a proof of every associated statement.

The specific baselines used here are as follows.

1. Gowers’ Lemmas 3.8–3.10 supply mixed cube inequalities, the norm triangle inequality, and the initial cube-count upper estimate. His Proposition 17.7 supplies the local phase-removal problem, with its derivative-energy input and cell-length scale.
2. The repository’s earlier centered-cube reports identify the quartic parallelogram coefficient, and Source 04 gives the sharp odd-order comparison at U^3 . We reuse those conclusions with attribution and include the quartic proof for clarity.
3. Sources 28 and 29 contain earlier norm-comparison bounds; Source 29 gives a central-binomial expression tending to $1/2$ and a lower construction tending to $1/3$. Source 29 is titled *Further Quantitative Refinements of Gowers’s Szemerédi Argument: higher norm comparisons, joint quadratic rigidity, and restriction by affine rank*. The present uniform limit answers that source collection’s remaining endpoint question.

4. Sources 34, 35, and 45 contain later fourth-order bounds. The comparison in this article uses the Source 45 coefficient $4795/832$ for the sixteenth-power comparison, equivalently $c_4^{\text{odd}} \leq (832/4795)^{1/4}$.
5. The separate 7 October 2026 manuscript *Local Phase Removal and Cubic Rigidity in Gowers' Szemerédi Framework* [259] provides the sliding-window and positive-cover interface and the denominator $\hat{D}_k$ used for comparison. Its underlying mechanism is reproved in Section IV.68; the dimension-dependent face measure and exact positive identity are the improvement here.

The archive's `SOURCE_AUDIT.md` records Git blob identifiers, file hashes, the separate predecessor's identity, and the normalization choices. Original source files are referenced, not redistributed with this contribution. The source comparisons are pinned to this inspection; subsequent repository integration can change the index and report numbering.

A targeted primary-source literature check also inspected Eisner–Tao on large Gowers–Host–Kra seminorms [7] and Christ–Iliopoulou on Riesz–Sobolev inequalities for compact connected groups [242]. The former concerns fixed-order extremizers and near-extremizers; the latter gives relevant rearrangement context for compact groups and interval models. Neither inspected statement supplies the uniform-in-group high-order comparison proved here. This is a description of the checked statements, not an exhaustive priority claim. The full search log is included in `review/literature_scope.md`.

V.238.2 Supplementary computation and its limits

The accompanying checks are designed to catch normalization, factorial, aliasing, and boundary mistakes. The mathematical theorems depend on their proofs, rather than on finite sampling. The archive includes runnable scripts and their result files.

Verification item	What is checked
Bounded correlation	The scalar inequalities, the averaging identity, and cyclic sharpness by exact rational arithmetic.
Fourier moments	1,296 exact moment and collision checks on six odd-order groups, including product groups.
Positive smoothing	273 exact multiplier checks, with 97 kernels containing actual character aliases.
Explicit norm rate	241 exact tail-specialization checks, 247 integer parameter checks, and high-precision evaluation of the finite comparison formulas.
Lattice face cover	10,980 target inequalities in 32 parameter cases, with exact weights and exact kernel arithmetic.
Figures	Plots generated from the stated formulas and recorded data, with a reproducible plotting script.

The displayed decimal evaluations of transcendental formulas use high-precision floating-point arithmetic. They are approximate values of proved expressions and are not separately certified interval bounds. Exact formulas take precedence over their rounded presentations. Finite rational checks are exact for the tested cases; they do not certify an unbounded range of parameters.

The principal norm proof and local face-cover argument were reviewed in separate reasoning passes. The included notes record checks of the distinct-pair coefficient, character aliases, signed cube products, the uniform quantifier order, the face inclusion–exclusion identity, phase polarization, the

powers of M, N, U , and the local geometric ceiling. These reviews found no unresolved gap in the stated arguments. They were carried out by ChatGPT instances and should be evaluated as internal mathematical review.

V.238.3 Suggested repository integration

The TeX source is supplied as a single article, with its figures, verification scripts, result data, and review notes. It can be compiled independently with `latexmk -pdf`. The existing paper transcription need not be edited to integrate this research contribution. A suitable research index entry should list the following claims separately:

1. the bounded centered U^2 inequality and sharp order- N correction;
2. the group-uniform limit $c_d^{\text{odd}} \rightarrow 1/3$, its finite parameter bound, explicit rate, and fixed-group corollary;
3. the asymptotically sharp quartic cube coefficient, with its retained higher-order remainder;
4. the exact positive-power face identity and lattice cover;
5. the finite localization coefficient and the asymptotic optimality $c_k^{\text{loc}} \sim v_k$.

For formalization, a useful dependency order is: scalar correlation constraint; finite averaging and sharp cyclic example; Fourier-tail lemma; positive smoothing; finite-spectrum recovery; uniform comparison; continuous face identity; lattice interpolation; sliding windows and polarization; local asymptotic assembly. The finite parameter statements should be preserved as interfaces, since they contain more directly reusable information than the limits alone. No Lean declaration or proof-assistant certificate is included or claimed by this article.

The integration notes should mark the limiting norm question as answered in the stated class, while retaining the fixed-order norm questions, bounded equality classification, exact cover problem, and second-order localization problem in the open-question list. This avoids converting an asymptotic result into a stronger fixed-parameter assertion during later consolidation.

V.239 Exact verification, implementation, and repository integration

Source 64, delivered Section 9.

V.239.1 What the included computation checks

The mathematical proofs do not depend on a search engine or an optimization solver. The accompanying `verify.py` is a second line of defense against normalization and transcription errors. It checks the sixty-four residues of the seed, the inner partition modulo 16, digit lifts through three layers, small exact cyclic grids for the equal-arc and nine-copy lemmas, the two rational slice constants, and the coefficient identities used in the realization argument. It also checks the short counterexample to the two-arc version and the valuation patterns in Corollary III.102.2.

All finite assertions use integers or `fractions.Fraction`. Decimal logarithms are used only for reporting the displayed exponents, not for deciding injectivity, a modular equality, or a theorem hypothesis. The output `certificate.json` records the seed images and the validation counts. The computations do not certify that the seed is optimal among all alphabets, or that a literature search is exhaustive.

V.239.2 A membership algorithm with exact arithmetic

The file `construction.py` implements digit alphabets and membership in the final sets, given a finite input coloring. For a supplied coloring $\phi : \{0, \dots, N-1\} \rightarrow \{0, \dots, r-1\}$, label set S , modulus m , and rational width $\alpha = u/v$, membership of $n \in \mathbb{Z}/P\mathbb{Z}$ is decided as follows. First compute

$$j = \left\lfloor \frac{9Nn}{P} \right\rfloor, \quad a = \left\lfloor \frac{j}{3N} \right\rfloor, \quad b = \left\lfloor \frac{j - 3aN}{3} \right\rfloor, \quad c = j - 3aN - 3b.$$

The color is $(3a + c)r + \phi(b)$. If its label is z and $y \equiv n^{k-2} \pmod{P}$ is in $\{0, \dots, P-1\}$, the membership condition is

$$0 \leq ym - zP, \quad v(ym - zP) < umP. \quad (\text{V.239.1})$$

Thus no floating-point evaluation of a torus phase is required. The expensive large-scale input is the finite coloring, not this membership test. The supplied implementation accepts that coloring as input; it does not claim to make the enormous theoretical thresholds computationally practical.

V.239.3 The formalization boundary

The inspected repository file `Combinatorics/Ramsey/Lean/GowersSzemerédi/Section04.lean` represents the Section 4 conjectures as non-asserting proposition-valued definitions. Its variable k in `conjecture_4_2` corresponds to progression length $k+2$; the present manuscript instead uses k for the progression length. The observed file blob is `987daf1643df5bd45a9e9202a08fb2f7c579a103`.

No existing formal theorem is replaced, and no proof-status entry is upgraded by this manuscript. A sensible first formalization target is the finite $(1, 5, 10)$ seed, followed by the general digit-lifting theorem and the all-moduli embedding estimate. These require only finite arithmetic and induction. The valuation obstruction can be a separate algebraic module. The circle transfer, Haar integrals, equidistribution and comparison with the repository's degree-uniformity predicate form a genuinely larger analytic chain. Corollary III.106.3 supplies the exact-density step that such a chain must not omit.

The recommended placement is a new research subdirectory, not the formal proof ledger:

```
Combinatorics/Ramsey/Research/GowersSzemerédi/
binomial-separation-six-progressions/
```

The accompanying `FORMALIZATION_PLAN.md` separates theorem statements, finite certificates, external inputs and still-unformalized analytic interfaces.

[write] **7 October 2026; sources 60–69.** The suggested subdirectory was not used: by its placement (39d3295c4) this source is source 64 of this report, with labels `gsr:fm:64:`, `gsr:cc:64:` or `gsr:fz:64:` and files under the prefix `64-binomial-separation-` (the formalization plan as `64-binomial-separation-FORMALIZATION_PLAN.md`).

V.240 Further research questions

Source 64, delivered Section 10.

Question V.240.1 (The six-binomial exponent). *Is $\beta_6 = 3$? More quantitatively, do there exist r -element $(1, 5, 10)$ -separating sets in integer intervals of length $r^{3+o(1)}$?*

The gap is now between 3 and $\log_4 80$ within the comparisons made here. The cyclic perfect-code obstruction means that searching for one exact r -symbol, r^3 -modulus seed cannot answer this question. Increasing-block, nonstationary, algebraic or genuinely nondigital constructions are natural alternatives.

Question V.240.2 (Optimal order and the endpoint eight). *Is there a constant C such that every r admits an r -element mirror-separating six-binomial set modulo some $m \leq Cr^3$?*

This is stronger than $\beta_6 = 3$ and would give $\rho_6(\alpha) \ll \alpha^8$ by Theorem III.105.1 without changing the coloring input. The distinction between an exponent statement and a uniform constant is essential at an endpoint.

Question V.240.3 (Cyclic capacity without perfect tilings). *What is*

$$\sup \left\{ \frac{\log |D|}{\log B} : D \subseteq \mathbb{Z}/B\mathbb{Z} \text{ is } (1, 5, 10)\text{-separating, } |D| > 1 \right\}?$$

Does it equal $1/3$, although the supremum is never attained by a perfect seed?

A productive finite search should output verifiable certificates and separate completed searches from time-limited failures. The valuation covering argument may help derive quantitative packing losses, but the present theorem only rules out equality.

Question V.240.4 (Improving the coloring rather than the labels). *Can symmetric six-progression-free colorings use $O(N^\gamma)$ colors for some $\gamma < 1/9$? Can the four-term exponent $1/4$ be improved?*

The exact scalar degree barrier shows where a direct replacement of the norm polynomial cannot work. A different degree regime, a nonlinear nonpolynomial map, or a multi-stage construction may escape it. Improving the coloring dimension directly improves the d/b term in the endpoint theorem.

Question V.240.5 (Beyond a fixed power). *Does $\rho_4(\alpha)$, or $\rho_6(\alpha)$, decay faster than every fixed power of α ?*

These are the superpolynomial-decay questions in the Ruzsa/DTZ program. Neither endpoint proved here resolves them. An endpoint power is a useful quantitative refinement but should not be confused with the full decay conjecture.

Question V.240.6 (Other coefficient patterns). *For which even lengths, or other symmetric point patterns, do the relevant positive coefficients admit small mirror-separating alphabets? What is the right replacement when non-mirror trivial partitions are unavoidable?*

The general CRT theorem supplies one family. At lengths with subset-sum collisions among the half-binomial coefficients, one must either control all allowed partition types geometrically or replace mirror separation by a more nuanced label condition.

Question V.240.7 (Quantitative analytic realization). *Can one turn the prime-modulus proof into practical bounds $P_0(\alpha, \eta, k)$ for a prescribed centered uniformity error η and progression-count error?*

The complete-sum estimate is explicit, but a full bound must also track the Fourier approximation to a rectangle union whose complexity depends on the chosen coloring and labels. Optimizing only the complete-sum exponent would not address that entire cost.

Question V.240.8 (Optimal transfer and formal verification). *Can a different circle lift improve on nine palette copies while retaining an exact or comparably strong progression measure? Can the full six-variable chain, including rational-density correction, be formally checked?*

The three-arc obstruction applies only to the particular auxiliary lemma. It is not an optimality theorem for all lifts. Formalizing the seed and the induction first would make future changes in alphabets inexpensive to certify; formalizing the analytic transfer would then turn quantitative input improvements into reusable theorems.

V.241 Conclusion

Source 64, delivered Section 11.

The research yields a concrete change in two quantitative resources: the six-variable label exponent drops from the generic 5 to $\log_4 80$, and the coloring-to-torus step can be made exact for mirror separation, with only nine palette copies. Keeping the prime family scale-uniform then gives $\rho_4(\alpha) \ll \alpha^5$ and $\rho_6(\alpha) \ll \alpha^{7.847232636916081\dots}$ without an epsilon loss. The proofs also identify two precise boundaries: perfect cyclic digit alphabets cannot reach exponent 3 for the six-binomial weights, and a single polynomial color of degree at most $k - 1$ cannot exceed dimension $k^2/4$. These boundaries guide, rather than close, the next stage of the program.

V.242 Explicit certificates and finite verification

Source 65, delivered Section 9.

V.242.1 A coefficient-level reconstruction algorithm

The algorithm in the accompanying code uses no sampling of f to recognize an obstruction. It starts with the symmetric basis coefficients of T . There are $\binom{n+d-1}{d}$ such bits, indexed by multisets of d basis indices. For $d \geq 4$ form a binary matrix D_T with n columns. Its rows are indexed by $i < j$ and multisets $\mathbf{a}$ of $d - 4$ basis indices, and its entries are

$$(D_T)_{(i,j,\mathbf{a}),k} = T(e_i, e_i, e_j, e_k, e_{\mathbf{a}}) + T(e_i, e_j, e_j, e_k, e_{\mathbf{a}}). \quad (\text{V.242.1})$$

For $d = 4$ the multiset is empty. Because the defect is alternating in its first pair and symmetric multilinear in its frozen slots, these coordinates determine its first-slot linear map completely. Thus

$$K_T = \ker D_T, \quad r(T) = \text{rank } D_T. \quad (\text{V.242.2})$$

There are $R = \binom{n}{2} \binom{n+d-5}{d-4}$ rows. Elementary row elimination uses $O(Rn^2)$ bit operations in the elementary unbatched model; this is a polynomial in the displayed matrix dimensions, not a polynomial in a compressed circuit representation of T .

If the rank is zero, the support coefficients c_A and (IV.52.5) give a primitive of T . If the rank is one, let u be the unique nonzero row of the row space. Choose a basis vector e_c with $u(e_c) = 1$ and compute

$$\Omega(a, b) = \Phi_T(e_c, \dots, e_c)(a, b), \quad v(x) = \Omega(x, e_c), \quad S = T - C_d^{u, v}. \quad (\text{V.242.3})$$

Canonical recognition proves that u, v are independent and S is integrable. The algorithm checks the latter conclusion independently by verifying that each basis coefficient of S depends only on its support. Formula (IV.52.5) then specifies a primitive of S from those support coefficients. A rank at least two gives the noncanonical energy cap without solving any continuous optimization problem. The formulas specify the recovery algorithm mathematically. The supplied Python function **recover** returns u, v, S ; the primitive regression routine checks the support-to-polynomial formula, rather than furnishing a polished symbolic pretty-printer.

V.242.2 Exact census and analytic regression checks

The verifier `code/verify.py` uses the Python standard library. It performs integer, binary-field, rational, and Gaussian-rational arithmetic only. It contains no numerical optimizer and no reliance on Python assertions that disappear under optimization.

The executed tensor census contains 38 044 cases: all symmetric binary tensors in degrees 3 through 10; all symmetric trilinear and quartic tensors in dimension three; and 172 seeded higher-dimensional or higher-degree cases, including explicit canonical tensors. For every applicable tensor, the checker compares independently computed notions: defect rank, integration by support coefficients, and the exact number of integrable contractions. Every radical-codimension-one case is reconstructed as canonical plus an integrable residual.

n	d	All tensors	Integrable	Canonical endpoint	Noncanonical
2	3	16	8	8	0
2	4	32	8	24	0
2	5	64	8	24	32
2	6	128	8	24	96
2	7	256	8	24	224
2	8	512	8	24	480
2	9	1024	8	24	992
2	10	2048	8	24	2016
3	3	1024	128	896	0
3	4	32768	128	2688	29952

For the last row, the noncanonical tensors split into 19 712 with $r(T) = 2$ and 10 240 with $r(T) = 3$. These exact counts agree with the class-count formula independently of energy optimization.

A contraction dynamic program also computes the exact energy at $f = 1$ for each tensor. Its maximum over nonintegrable tensors matches b_d in every exhaustive family, and each computed value lies below the relevant radical-sensitive cap. This is a check on the theorem and its sharpness examples, not a computation of $\mathcal{M}(T)$ for every tensor.

There are also 34 tests on exact Gaussian-rational functions, including zeros and non-unit amplitudes, checking energy bounds and exact slicing. Four independently accumulated labelled-cube averages are compared with selected Fourier-square averages. Another 360 checks

verify derivative conjugation periods, and 18 test the constrained pure-tensor lemma. The primitive formula is checked on 480 direction tuples at all base points, with the next derivative checked in each case, in dimensions one through three and degrees two through six. Finally both induction recurrences are verified as rational identities through degree 1000.

V.242.3 What verification does not mean

The finite computations cannot establish an inequality for every complex function or every dimension. The full proofs above establish those claims. Neither a finite phase grid nor the constant-function census is substituted for the continuous upper-bound argument. The code is not a Lean formalization, and the manuscript has not undergone independent refereeing. The proposed formalization plan does not change any checked statement count or theorem-status ledger in the repository.

V.243 Relation to Gowers’s argument and the formal interfaces

Source 65, delivered Section 10.

V.243.1 The precise phase-removal connection

The relevant stage of Gowers’s paper is Section 17, especially Lemma 17.1 and Proposition 17.2 [1]. That stage converts structured derivative frequencies into polynomial phase information. The present result is a characteristic-two counterpart for a supplied, fully symmetric homogeneous tensor, where factorial division does not supply a primitive. It quantifies precisely the high-energy obstruction to integration rather than assuming symmetry suffices.

In the formal statement inspected at repository snapshot `8c40adc24df2`, `Sections17_18.lean` makes a prime cyclic modulus and invertibility of $(k + 1)!$ explicit in `proposition_17_2`; it records degree at most $k + 1$ for the phase, correcting the printed degree. The selected-frequency normalization here corresponds to $d = k + 1$: divide its unnormalized Fourier-square sum by $|V|^{d+1}$. One must additionally reconcile the direction/conjugation convention: the repository uses $f(s)\overline{f(s - a)}$, while this article uses $f(x + h)\overline{f(x)}$. On a Boolean group these differ by conjugation at each differentiation; the selected squared Fourier magnitudes at Boolean characters agree. This is an interface check, not a silent identification of the ambient groups.

V.243.2 Three limitations that matter mathematically

First, the result concerns $\mathbb{F}_2^n$, not arbitrary cyclic groups, Bohr neighborhoods, or progression cells. Translation invariance of the full uniform measure is used in the Fourier cancellations and slicing. Second, the primitive is circle-valued and can have genuine dyadic depth; it is not required to be a classical polynomial over the base field. Third, the sharp bound assumes a supplied fully symmetric tensor. It does not symmetrize arbitrary multilinear data, handle unrestricted mixed vertex functions, or produce a selector at low energy.

The inspected `Definitions.lean` defines `IsMultilinear` using all square-free monomials, including the constant term. In standard terminology this is multiaffine, unlike the homogeneous tensor convention used here. Removing those extra terms would require a separate argument. Thus

the theorem is not presented as a drop-in replacement under identical hypotheses for the prime-cyclic proposition.

An improved quantitative Szemerédi estimate would require a complete selector-construction and localization argument, followed by all remaining density-increment losses. No such global estimate is inferred from the local extremal theorem.

V.243.3 Provenance and formalization boundaries

The inspected repository snapshot is

8c40adc24df2e6f338cea728887772aead4bf144.

The interface blobs actually read were

```
Sections17_18.lean:
18495f50475835ca1c497d85b0052215f865b8ca
Definitions.lean:
97113b7afa6925a2dd4b76641eeaeff09597ab6a.
```

The research README and Boolean-phase formalization notes were inspected separately. The decisive newer comparison manuscript [264] was read in full from the user’s Library; its own comparison snapshot is 5c9a442d263632fdcd4e9690154c12c2dcc70b53. We do not conflate that manuscript’s snapshot with the one used here.

The clean formalization dependency order is: finite cube identities; Boolean integration by basis coefficients; cubic projective orthogonality; the canonical theorem; defect-radical recognition and the integrable-slice identity; and finally the all-degree induction. Once those ingredients are present, the genuinely new induction requires only finite averaging and elementary inequalities. The companion `FORMALIZATION.md` records suggested declarations and explicit hypotheses. No research-result axiom, unproved placeholder, or repository build result is included as evidence of verification.

V.244 Further research questions

Source 65, delivered Section 11.

These questions follow the proved results. The septic threshold, all-degree universal value, and endpoint tensor classification are no longer listed as open questions within this manuscript.

Research question V.244.1 (The second extremal value). For fixed $d \geq 4$ and n , determine the maximum of $\mathcal{M}(T)$ over the noncanonical, nonintegrable tensors, when that family is nonempty. Is the cap $1 - 3d/2^{d+1}$ ever attained? Which extra loss is forced because optimal contraction functions must all arise as derivatives of one common function? The retained profile (IV.58.7) is a first computable refinement, not a solution of this optimization problem.

Research question V.244.2 (Rank-sensitive extremizers). Determine the exact extremum when the defect radical has prescribed codimension $r \geq 2$. Is the distribution of further contraction radical ranks enough to determine it, or do tensors with identical contraction profiles have different maxima? Constructing such a pair would locate a precise limitation of profile-only arguments.

Research question V.244.3 (Maximizing functions and phase stability). Classify all f with $\mathcal{E}_{C_d}(f) = b_d$ and all near-extremizers modulo natural phase symmetries. The tensor classification and (IV.57.5) do not reconstruct phases. Can the canonical slice-deficit conditions yield a dimension-independent distance estimate to a fully described extremizer family? Any proposed uniqueness statement must allow multiplication by phases with zero d th derivative.

Research question V.244.4 (Inactive coordinates outside the canonical class). For which tensors T_0 does a surjection $q : V \rightarrow V_0$ satisfy $\mathcal{M}(q^*T_0) = \mathcal{M}(T_0)$? The canonical case holds in all degrees, and the cubic rank law gives it for every cubic tensor. The noncanonical higher-degree problem remains. A tensor’s dependence on few coordinates alone does not constrain an arbitrary optimizing function to those coordinates.

Research question V.244.5 (Direct sums and compatible slice optimization). Determine whether maxima multiply under direct sums of canonical tensors in degrees at least four. Product functions give the lower bound, while the cubic law proves equality in degree three. Can non-product functions improve the energy in higher degrees? More generally, characterize when the upper bound in exact-slicing relaxations is attainable by one function.

Research question V.244.6 (Odd characteristic). Develop an analogue of the integrable-contraction radical for the p -repetition defect in characteristic $p > 2$. Does a minimal-radical classification isolate the exceptional obstruction family and let an all-degree induction close? The Boolean derivative-period lemma is not valid unchanged in odd characteristic, so both the canonical family and its constrained energy must be rederived.

Research question V.244.7 (Local measures and boundary errors). Quantify the slice and Fourier-period errors for weighted averages, Bohr neighborhoods, and progression cells. What translation-invariance hypotheses preserve a useful version of the rigidity gap γ_d ? A successful local theorem must compare boundary losses with an exponentially small degree-dependent margin rather than hide them in an unspecified error term.

Research question V.244.8 (Multiaffine and incompletely symmetric selectors). Give a quantitative reduction from the repository’s multiaffine selector interface to genuinely homogeneous symmetric tensors without discarding lower-order terms or silently changing the energy. At which cost can the sharp certificate here be inserted after such a reduction? This question is distinct from certifying a tensor already known to be symmetric.

Research question V.244.9 (Certified reconstruction and formal proof). Implement the matrix certificate, primitive construction, and endpoint recognition in Lean, then formalize the cubic Hilbert-space estimate and the all-degree induction. Can a small verified checker accept explicit row-reduction and phase certificates independently of the search procedure that found them? The input representation and the distinction between a proved theorem and a proposition-valued catalogue entry should remain explicit throughout.

V.245 Conclusion

Source 65, delivered Section 12.

The universal extremal problem is governed by the same canonical obstruction in every degree. The missing ingredient was not another higher-order support census: it was the exact identification of the integrable slice set with a linear radical. Codimension one gives the canonical family, and larger

codimension forces a strict loss under slicing. This closes the universal induction and excludes every noncanonical endpoint.

The resulting theorem is sharp over all bounded complex functions and all finite Boolean dimensions. Its geometric consequences are also exact: endpoint tensors are classified modulo integration, their classes are counted, and all maximal integrable restrictions are identified. The remaining work lies in phase stability, noncanonical optimization, and the localization and selector interfaces needed for broader quantitative applications.

V.246 A small explicit certificate

Source 65, delivered Appendix A.

In dimension two, let $u(x) = x_1$, $v(x) = x_2$. For the quintic $C_5^{u,v}$, the symmetric basis coefficients, listed by the number of occurrences of e_2 from zero through five, are

$$(0, 1, 0, 0, 0, 0).$$

The defect is $u \otimes u$ in its two frozen slots, multiplied by $u \wedge v$. The first-slot matrix has row space spanned by u , so $r(C_5) = 1$ and its maximum is $13/16$. Changing the tensor by any of the eight integrable binary quintics leaves its maximum unchanged.

For contrast take $T = C_5^{u,v} + C_5^{v,u}$, with coefficients

$$(0, 1, 0, 0, 1, 0).$$

Its defect is $(u \otimes u + v \otimes v)(u \wedge v)$, whose first-slot radical is zero in dimension two. Thus $r(T) = 2$. Exactly one quarter of its contractions are integrable and its energy is at most $49/64$ for every bounded complex function. In particular it cannot attain the universal endpoint $13/16$. The assertion follows from the general proof, not from checking a finite collection of phases.

V.247 Reproduction

Source 65, delivered Appendix B.

From the package directory, run

```
python3 code/verify.py
python3 -O code/verify.py --output data/verification_optimized.json
latexmk -pdf -interaction=nonstopmode -halt-on-error \
    all_degree_boolean_integration.tex
```

The two JSON verification records must agree. The `Makefile` provides the equivalent targets `check`, `pdf`, and `all`. The PDF is compiled from the supplied TeX source. The package includes source and proof audits and a formalization plan, but no prior manuscripts, third-party paper PDFs, or font files.

V.248 Exact recovery algorithms and reproducibility

Source 66, delivered Section 12.

V.248.1 Recovering a partition from exact input

For a positive ordered set C represented by exact rational numbers, the following procedure returns the normal form whenever its hypotheses hold. First compute $\sigma(C)$ by testing each positive difference against the set, and put $r = \sigma(C)$. If $r = 0$, return the arithmetic progression. If $n \geq 3r + 1$, divide every point by c_1 and form

$$\lambda_i = c_i/c_1 - i.$$

Theorem II.302.1 guarantees that these are nonnegative, nondecreasing integers with total r . No rational approximation, integer-relation heuristic, or tolerance test is needed. For a general exact ordered field, the same procedure is valid provided its ordering, arithmetic, and equality operations are available.

For an endpoint set A , use $d = \max A - \max(A \setminus \{\max A\})$, translate its minimum to zero, and divide by d . Under the theorem's hypothesis all coordinates become integers, the missing positions give H , and (II.303.1) verifies the certificate. The output includes both a structural certificate and an exact value of the energy via (II.304.5). The theorem does not require enumeration of all partitions in order to recover the particular input's partition.

With constant-time exact arithmetic and hashing as an abstract cost model, the direct defect computation takes $O(n^2)$ membership tests and $O(n)$ storage for the set. Rational bit complexity also includes the cost of exact arithmetic on the input numerators and denominators. Without hashing, binary search in the sorted set gives $O(n^2 \log n)$ comparisons. Enumerating all admissible templates takes $p(r)$ outputs, and its cost is independent of the ambient coordinate sizes except for writing the templates.

For the four-level energy test, compute the difference multiplicities, obtain $E(A)$, and compare it to the exact integer threshold. If the threshold holds, the sorted set can be compared against the eight normalized templates. A hypothetical discrepancy is a failed theorem certificate, not something the algorithm should conceal by rounding.

V.248.2 What the supplied program checks

The file `code/verify.py` is a dependency-free Python program using integers and exact rational arithmetic. Its assertions remain active under Python's optimized execution mode because failures are raised explicitly. Its default finite checks are as follows.

Among the 65,535 positive sets, 108 satisfy the proved normal-form criterion, including zero-defect cases. Among the endpoint sets, 58 satisfy the positive-defect envelope criterion. These smaller numbers are reported to distinguish the number of inputs examined from the number lying in the narrow theorem hypotheses. The recorded output contains no counterexample to any proved statement.

The program also tests the conjectural threshold $n \geq 2r + 2$ on the same finite positive-set box. There are 103 applicable positive-defect instances and no counterexample in that box. This diagnostic is explicitly separated from the proved theorem, and does not justify replacing the theorem's cutoff in an application.

The file `data/partition_spectrum.csv` contains every boundary hole set of weight at most 12 and its exact κ value. `data/verification.json` records parameters and counts, and `data/verification.log` gives the human-readable run summary. The formulas are proved before

Exact finite test family	Instances
All nonempty $C \subset [1, 16]_{\mathbb{Z}}$; Schur defect and applicable normal-form conclusions	65,535
All $A \subset [0, 16]_{\mathbb{Z}}$ containing zero with $ A \geq 2$; energy identities, recurrences, applicable gap and equality tests	65,535
Partition templates of weights $1 \leq r \leq 12$	271
Boundary-hole energy identities at three interval lengths per template	813
Explicit threshold obstructions, $1 \leq r \leq 40$	40
Eight energy models, $5 \leq m \leq 70$	528
Exact rational sets from a fixed deterministic seed	700
Rationally dilated partition models, $1 \leq r \leq 8$	66
Rank-one embeddings of the eight models in $\mathbb{Z}^3$	112

Table V.248.1: Reproduced checks in the accompanying archive. The exhaustive boxes are finite and do not certify the theorems beyond those boxes.

these computations are described; none of the proofs relies on a floating-point calculation or a finite search standing in for an infinite argument.

V.249 Further research questions

Source 66, delivered Section 13.

The questions below are ordered approximately by their proximity to the proved mechanisms. They are not claims that a global inverse theorem or a Ramsey bound has already been improved.

Conjecture V.249.1 (The optimal Schur threshold). For every integer $r \geq 1$, Theorem II.302.1 remains valid under $n \geq 2r + 2$.

Proposition II.302.4 would make this threshold optimal. The current proof obtains an initial arithmetic block of length $n - 2r$ and requires this block to exceed the entire defect budget. An improvement must exploit the cost of the tail more efficiently, rather than simply repeating that comparison. The first task is to classify the possible cancellations from sums of two missing points.

Question V.249.2 (Critical-size exceptional families). *Classify all positive sets with $n = 2r + 1$ and defect r that are not partition normal forms. Is the central-hole family of Proposition II.302.4 the only large- r mechanism up to positive dilation, or can several interacting holes realize the same critical ratio?*

This question asks for an equality theory at the obstruction threshold, not just an improved sufficient bound.

Conjecture V.249.3 (The second boundary-energy constant). For $r \geq 3$, among hole sets of weight r other than $\{r\}$,

$$\max \kappa(H) = r^2 - 2r + 4,$$

with equality only for $H = \{1, r\}$.

The displayed candidate has that value by direct substitution into (II.304.3). The partition data through weight 12 satisfy the conjecture. A proof would give the next sharp endpoint-energy gap at every defect, extending the all- r envelope beyond its first equality case. The crude bound $\kappa \leq r \max H$ is not enough: it loses the interaction between the number of holes and the pair-overlap term.

Question V.249.4 (The weighted partition spectrum). *Determine the bivariate generating function, or a useful recurrence, for*

$$F(z, w) = \sum_{H \subset \mathbb{Z}_{>0} \text{ finite}} z^{\text{wt}(H)} w^{\kappa(H)}.$$

Which distinct partitions have the same κ ? Can the number of energy values at fixed weight be bounded substantially below $p(r)$?

The two weight-four examples in Table II.304.1 show that injectivity already fails at low weight. A useful recurrence would separate those combinatorial coincidences from accidental equalities between energies at different weights and cardinalities.

Question V.249.5 (The fifth and later global energy levels). *Find the fifth largest energy, with all equality cases, for sufficiently large m , and then develop a systematic method for any fixed number of levels. The one-hole model $[0, m]_{\mathbb{Z}} \setminus \{3\}$ has $D = 3m - 12$, suggesting a next scale, but this alone supplies no lower bound for arbitrary sets outside the first four levels.*

The successful third- and fourth-level proofs distinguish arithmetic, second-level, and other prefixes, then reduce the last case to a small endpoint budget. A general method will require a controlled list of prefix types with holes near *both* endpoints. New mixed-end models should be expected; the class C_m is the first warning.

Question V.249.6 (A full inverse theorem at linear total defect). *For fixed K , does $D(A) \leq K|A|$ force, for sufficiently large $|A|$, a bounded-hole one-dimensional description of all but a bounded number of points, with constants depending only on K ? Can a stronger whole-set statement hold after listing a finite family of exceptional endpoint patterns?*

Theorem II.302.1 has a small *local* defect assumption, whereas this question controls the sum of local defects. Bridging that quantifier difference is the main obstacle. An arbitrary low-defect prefix cannot be selected without controlling how the omitted endpoints interact with it.

Question V.249.7 (Robustness to perturbed relations). *Develop a quantitative analogue when additive equalities are known only up to an error tolerance, or when a small proportion of their labels are corrupted. What separation assumptions on the real set are necessary to recover a nearby arithmetic template?*

Exact order and exact reflection are decisive in Lemma II.301.1. In a metric version, accumulated reflection errors may be amplified by a long successor chain. A useful result must specify both an additive error scale and an arithmetic separation scale; an unspecified “approximately equal” version would be false or vacuous after rescaling.

Question V.249.8 (Weighted and multiplicity-sensitive extensions). *Can the reflection method yield an exact or stable classification for nonnegative weights, or for multisets, near the corresponding energy maximum? Is there a weighted substitute for the partition budget that retains sharp endpoint information?*

The current partition count uses integrality of each failed pair and strict ordering of distinct points. Those are not cosmetic features: weights can split a unit defect continuously, and multiplicities change the maximum itself.

Question V.249.9 (Rectification of near-extremal cyclic graphs). *Identify hypotheses under which a cyclic graph of very high selected energy admits a torsion-free lift with the same additive quadruples, and track the resulting losses explicitly. Can the affine conclusion of Corollary II.309.3 then replace a more expensive local inverse step in a restricted branch of Gowers’s argument?*

Proposition II.309.4 is a sufficient interface, not a way to produce the needed short lift. Subgroup examples rule out a universal transfer with no additional assumptions.

Question V.249.10 (Formal certification and effective template generation). *Formalize the all- r normal form and the four-level classification, then implement a verified certificate generator returning the partition or a failed hypothesis. Can the finite template recursion from Question V.249.5 be made proof-producing rather than merely experimental?*

The reflection lemma, the no-two-holes count, and the finite prefix case splits are natural modular targets. A formalization should prove the integer-valued defect identity before using integrality to exclude an intermediate energy level.

V.250 Conclusions

Source 66, delivered Section 14.

The ordered-energy maximum used in Gowers’s Section 7 admits a discrete refinement with a particularly rigid local structure. Once the final Schur defect is smaller than a fixed fraction of the set size, all possible configurations are indexed by integer partitions. Their energy is explicitly computable from boundary holes, and the best energy at each endpoint defect has a unique one-hole normal form. The low-defect instances complete two equality analyses beyond the second energy level and yield full-set recovery in the first-four-level window.

These conclusions are exact finite theorems, not asymptotic guesses. Their limits are equally specific: the sufficient Schur threshold is not proved optimal; the global fifth level is not determined; the linear-defect inverse problem is not solved; and the passage to the cyclic graphs and density iterations in the general Szemerédi proof requires additional work. The partition theorem and the explicit energy constants give concrete tools with which to pursue those next steps.

V.251 The relation of the first gap to inverse Pollard theory

Source 66, delivered Appendix A.

This appendix is an attribution check, not an input to any proof above. For a finite integer set A of size m , let

$$r_A(t) = \#\{(a, b) \in A^2 : a - b = t\}, \quad S_j(A) = \sum_t \min(j, r_A(t)).$$

Embedding the finite set in a sufficiently large prime cyclic group avoids all wraparound. Pollard’s inequality gives

$$S_j(A) \geq j(2m - j).$$

For a non-arithmetic-progression set, the equality classification of Nazarewicz, O'Brien, O'Neill and Staples [190, Theorem 3], together with the $j = 1$ equality case recalled there, makes this inequality strict for $1 \leq j \leq m - 2$. The exceptional alternative involving sets of size $j + 1$ is excluded by $j \leq m - 2$, and the large-modulus alternatives are excluded by choosing the prime sufficiently large. Since $r_A(t) = r_A(-t)$ and $r_A(0) = m$, one has $S_j(A) \equiv j \pmod{2}$. The arithmetic-progression value $j(2m - j)$ has the same parity. Strictness therefore gains at least two for every $1 \leq j \leq m - 2$. The scalar identity, valid for $0 \leq u \leq m$,

$$u^2 = (2m - 1)u - 2 \sum_{j=1}^{m-1} \min(j, u),$$

yields

$$E(A) = (2m - 1)m^2 - 2 \sum_{j=1}^{m-1} S_j(A).$$

Comparison with an arithmetic progression gives

$$E(A) \leq M_m - 4(m - 2).$$

Thus the first gap is already encoded in published inverse Pollard theory. Its elementary proof in Section II.306 should not be described as a new extremal inequality.

For a finite subset of a torsion-free abelian group, this particular attribution argument can first use a Freiman 2-isomorphism to the integers: identify its generated subgroup with $\mathbb{Z}^d$, and choose an integer linear functional nonzero on every nonzero vector in the finite set $2A - 2A$. Such a functional exists because finitely many proper rational hyperplanes cannot cover $\mathbb{Z}^d$. This preserves all quadruple relations. It is a finite-set construction, not an injective group homomorphism from $\mathbb{Z}^d$ to $\mathbb{Z}$ when $d > 1$.

V.252 Proof dependencies and formalization plan

Source 66, delivered Appendix B.

The article can be formalized independently of analytic uniformity norms. A suitable new namespace is `GowersSzemerédi.OrderedEnergy`; this is a proposal, not a claim that such a module already exists in the repository.

The main dependency chain is

$$\begin{aligned} &\text{endpoint energy recurrence} \implies \text{ordered-defect identity,} \\ &\text{two zero defects} \implies \text{long initial arithmetic block} \implies \text{partition normal form,} \\ &\text{partition normal form} \implies \text{boundary-hole models} \implies \text{sharp endpoint envelope,} \\ &\text{defects 0, 1, 2, 3 and prefix case splits} \implies \text{third and fourth equality classifications.} \end{aligned}$$

An implementation can begin with a strictly increasing finite sequence of real numbers rather than repeatedly sorting a finite set. The local defects e_i and endpoint defect should be natural numbers. Define the total defect first by the ordered triple count, prove the energy identity $E(A) + 4D(A) = M_m$, and only then identify it with the quotient used in this paper. That order avoids division and nonnegativity obligations at every induction step.

Proposed milestones are as follows.

- F1. `energy_endpoint_recurrence` and `ordered_defect_identity`: finite bijections and a polynomial identity in m .
- F2. `schur_two_zero_prefix`: two finite-set reflections and the successor-chain argument. This is the main order-theoretic lemma.
- F3. `schur_partition_normal_form`: a first-failure argument for lattice membership, then a first-failure argument for the $2L + 1$ bound, followed by inclusion–exclusion.
- F4. `endpoint_hole_normal_form` and `boundary_hole_energy`: reflection, sums of integer intervals, and the explicitly counted overlap term.
- F5. `kappa_le_weight_sq` with its equality case; `endpoint_sharp_energy_bound`.
- F6. `third_energy_classification` and `fourth_energy_classification`: induction for the earlier gap, then exhaustive mathematical case splits on the prefix type and the small endpoint defect.

No source file in this package claims these declarations have been compiled. The existing proof-status ledger should not be changed to “proved” on the basis of this manuscript or its Python checks. The appropriate integration is a research article plus proof targets; a later verified development can promote the corresponding statements separately.

V.253 Source audit and reproduction notes

Source 66, delivered Appendix C.

The original-paper comparison uses the repository transcription of [1], specifically the energy paragraph immediately before Theorem 7.2 and the subsequent Proposition 7.3. The observed source blob is

41085a0efde8c1932e86e80c791984f52841c22e.

The prior-result comparison uses the source-45 energy review [266], observed blob

7984f49e51f4cfb0e9cd44f8385b5302bc75242f.

The local-refinements README was also inspected, with observed blob

1ec635141ba5b9fb2ceff2a4a1ac59107de9587b.

These are content identifiers for the files actually inspected; the article does not assert that every repository file was surveyed. Full paths and source URLs are recorded in `SOURCES.md`.

The published inverse Pollard paper was read at its publisher’s PDF, including visual inspection of Theorem 3 on printed page 2. Its exceptional set-size alternative was checked explicitly. The bibliographic information for Gowers’s paper was checked against the publisher page. A limited web search did not resolve the priority of the all-defect classification or of the complete third- and fourth-level equality statements. That limitation is why priority is left open.

To reproduce the finite checks from the package root, run

```
python3 code/verify.py
```

To compile the paper, run

```
pdflatex -interaction=nonstopmode -halt-on-error article.tex
pdflatex -interaction=nonstopmode -halt-on-error article.tex
pdflatex -interaction=nonstopmode -halt-on-error article.tex
```

The supplied `Makefile` also exposes these commands. The mathematical verification needs only Python's standard library; the PDF build needs a LaTeX installation with the packages listed in the preamble. The PDF was compiled and visually inspected during package preparation. The document is a conventional mathematical proof draft, not a certified theorem-prover output.

V.254 Further research: precise questions and plausible routes

Source 67, delivered Section 7.

The results above replace several vague improvement targets by well-defined extremal problems. The following questions are not assertions that an answer is already known to follow from our proofs. The first group asks for completion of an exact classification; the second asks for structural information that the current local interfaces discard; the last asks which improvements can survive the full Gowers iteration.

Here $d(f, g) = |G|^{-1} |\{x \in G : f(x) \neq g(x)\}|$ denotes normalized Hamming distance, and $d(f, \text{Aff}) = \min_{\ell \text{ affine}} d(f, \ell)$.

V.254.1 The global unequal-budget Fourier problem

Question V.254.1 (Phase regions beyond the strict matrix branch). *For fixed R and a , determine $V_{R,a}(\varepsilon)$ throughout the nonnegative budget orthant. Classify the regions on which its extremizing support and active Fourier constraints have constant type.*

Theorem I.76.2 supplies an exact open region and a certificate that can be checked beyond the small neighbourhood used to prove its existence. Its boundary can occur when a Fourier coefficient of the certificate reaches zero, a moment matrix loses rank, an atom loses mass, or the arc mass reaches zero. These events have different meanings. A vanishing certificate coefficient permits an inactive moment constraint; a vanishing atom changes the support combinatorics. Treating all such events as the failure of one numerical solver loses this information.

A concrete first project is the complete $R = 3$ classification for an algebraic value of $\cos a$, beginning with $a = \pi/3$. On the current region the answer is rational in the budgets. The next task is to derive adjacent certificates and prove their regions meet without leaving a gap. General semidefinite optimization supplies a useful search mechanism, but a solution should give exact formulas or exact algebraic descriptions, including boundary multiplicities and the open-arc convention.

Question V.254.2 (The next finite-grid term). *For fixed R, a in the first window, determine the subsequential limits of*

$$N^2(\rho_{R,N}(a) - \rho_R(a_N)), \quad a_N = \frac{2\pi}{N} \left\lceil \frac{aN}{2\pi} \right\rceil,$$

Here $\rho_{R,N}(a)$ denotes the infimum of $\max_{1 \leq j \leq R} |\hat{\mu}(j)|$ over probability measures supported on the grid $2\pi\mathbb{Z}/N$ outside I_a , as in Theorem I.76.8.

The $O(N^{-2})$ error is now proved; the remaining issue is its exact coefficient and possible oscillation. The nonendpoint atoms of the continuous extremizer have their own positions relative to the grid. A second-order answer is therefore likely to involve several rounding phases. A useful result would separate the geometric endpoint drift, the barycentric rounding error at each atom, and reoptimization of the atomic weights. It should also decide when the coefficient $\pi^2 R^2/2$ in the universal upper bound can be improved uniformly. The limiting contact geometry from the exact Fourier analysis gives a starting point. The comparison measure has zero arc mass, at the boundary of the strict matrix region; a separate boundary sensitivity analysis is therefore needed. The interior Hessian theorem cannot simply be applied at that endpoint.

V.254.2 Beyond ternary target geometry

Question V.254.3 (The largest target class with the full cubic range). *For which abelian targets W does*

$$1 - \mathcal{E}(f)/|G|^3 \geq \Psi(d(f, \ell)) \quad \text{whenever } d(f, \ell) \leq 2/3$$

hold for every finite abelian domain G , map $f : G \rightarrow W$, and affine reference ℓ ? Determine the equality cases for each such target class.

The new result answers this for exponent-three targets. The earlier manuscript proves a smaller sufficient range for targets without involutions, and a further improvement when order-three elements are also absent [271]. The ternary proof cannot simply be copied: its exact projection factor and seven target configurations use the specific dual geometry of $\mathbb{F}_3$.

Two realistic first steps are cyclic targets of orders five and seven, and targets of exponent p for a fixed odd prime. One can compute the corresponding projected kernels and retain all relation types before taking a pointwise maximum. A counterexample would also be valuable, but should identify which relation pattern defeats the cubic and whether the offending density can be moved into the range relevant to a nearest-affine inverse theorem.

Question V.254.4 (Sharp stability near the coset extremizers). *For t in a compact subinterval of $(0, 2/3)$ and $\varepsilon(f) \leq \Psi(t) + \eta$, determine the best quantitative distance from $f - \ell$ to a one-color subgroup-coset error. Track the dependence on t and the discrete set of possible subgroup indices.*

Corollary II.294.7 already bounds the mass outside a dominant error color and the proportion of target triples with rank at least two. The remaining step is spatial: a nearly extremal support must be close to a subgroup coset. The one-color identity expresses the exact excess in terms of $s^3 - E(S)$, so quantitative stability of the maximal set-energy inequality is the natural input. A useful theorem should combine color and support stability without losing the correct dependence as t approaches either endpoint. Near $t = 2/3$, the additional affine equality family must be incorporated instead of forced into the interior coset description.

Question V.254.5 (The attainable affine-distance/energy spectrum). *For fixed G and an exponent-three target, determine the possible pairs*

$$\left(d(f, \text{Aff}), 1 - \mathcal{E}(f)/|G|^3 \right),$$

and characterize the lower boundary. Seek uniform descriptions for families of cyclic groups and elementary abelian 3-groups.

The two recovery thresholds and the increasing inverse branch are exact in their stated senses. They do not identify the lower envelope at every real parameter. Equality with Ψ requires a subgroup coset, so the allowed equality densities are arithmetic, not continuous. The sharp non-affine maximum identifies the first nonzero point when $N \geq 13$; the next support size and the next attainable defect remain meaningful finite problems. Exhaustive enumeration can supply conjectures, but results should distinguish fixed-domain spectra from limits obtained by changing the group.

Question V.254.6 (Sample complexity at the exact recovery thresholds). *Determine the optimal query and sample complexity for recovery below graph defect $1/4$, or below $4/9$ on a ternary domain. Obtain near-linear running time where possible, with explicit dependence on the margin from the threshold and on erasures.*

For an explicit table, a polynomial-time procedure is already immediate: enumerate the at most $|G|$ homomorphisms $G \rightarrow \mathbb{F}_3$ and their three affine translates, and find a nearest scalar affine map in each target coordinate. The global theorem guarantees that these are the coordinates of its unique vector-valued affine map, so combining them retains its sharp distance bound. The proof using all projections establishes that compatibility; the algorithm need not inspect them all. The open issue is sharper resource complexity, especially for oracle access and partial observations. For partial graphs, completion independence is already proved. A useful sampling theorem should exploit it without filling in all missing values, and should state whether erasures are random or adversarial.

V.254.3 Information that permits interpolation compression

Question V.254.7 (Surplus anchor agreement and compatible labels). *Suppose each relevant polynomial class contains $u > t$ valid anchors. What additional compatibility hypothesis gives a sharp improvement over $E_t(L_1, \dots, L_r)$, and what is the optimal dependence on u and the list-size profile?*

The exact lower construction uses each target polynomial at precisely t selected anchors. A genuine surplus can therefore defeat that construction. However, dividing the count by $\binom{u}{t}$ is justified only when all those anchor representations describe one common candidate on the entire base cell. Pointwise choices that depend on the base point need not do so. A first theorem should state the minimal compatibility condition under which multiplicity counting is valid, then test whether Gowers' product property implies that condition, perhaps after a controlled refinement.

The finite-field and partition variants are also concrete. When the field is too small for the generic construction, collisions can reduce the number of distinct interpolants. If every further cell must have a prescribed arithmetic width, the present accounting $bQ \geq E_t$ is not the final geometric constraint. These questions are described in more detail at the end of Section II.339; they require additional incidence or product information, rather than another count of the same t -subsets.

V.254.4 Completing and escaping the norm-layer barrier

Question V.254.8 (Critical odd norms for every progression length). *For every $m \geq 2$ and prime $2m < p < (2m - 1)^2$, is*

$$M_{2m+1,m}(p) > 0?$$

Equivalently for the classification problem, must the norm of degree $2m + 1$ contain a nontrivial symmetric $2m$ -term progression?

All larger odd degrees are already excluded over every allowed characteristic. The exceptional fields are prime fields, since a proper power of a prime larger than $2m$ already exceeds $(2m - 1)^2$. Thus there is a finite and explicitly bounded set of checks for each length. Lengths four, six, and eight are settled in this article.

A structural approach seeks an irreducible polynomial of degree $2m + 1$ with odd part

$$X \prod_{i=1}^m (X^2 - (2i - 1)^2).$$

Its root yields the required norm equalities. Taking only a constant even part is sufficient in the eleven eight-term examples, but it would be premature to require that stronger ansatz for every field and length. Allowing an arbitrary even part provides more freedom. An existence theorem for this prescribed odd part would remove the last condition from the independent norm-block classification.

Question V.254.9 (Unrestricted scalar quintics for four-term symmetry). *What is the largest n for which a scalar polynomial $P : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$ of total degree at most five can avoid every nontrivial symmetrically coloured four-term progression, over fields of characteristic greater than five?*

The repository's general degree filtration leaves a gap between the known dimension-four construction and the general upper bound eight [272]. Theorem III.109.1 shows that independent norm blocks of any degree cannot close that gap by producing a larger example. A successful construction would need interacting terms or a different algebraic design. A sharper impossibility theorem would need to use relations between the coefficient forms at different centres, rather than treating the pointwise equations as an arbitrary system of their stated degrees.

One practical programme is to classify small-rank leading quintic forms up to linear change of variables, then examine the induced cubic and linear difference coefficients. A purely random search over all coefficients is unlikely to explain the mechanism. The target should be either a reproducible explicit example in dimension at least five, or a structural theorem ruling out a well-defined larger class than independent norms.

Question V.254.10 (Interactions and quantitative supersaturation). *How robust is the norm-layer obstruction under cross-block perturbations? Determine assumptions on a polynomial interaction $R(x_1, \dots, x_s)$ that force a positive lower bound for the number of symmetric progressions in $\sum_j v_j N_j(x_j) + R$.*

An independent extra summand cannot repair a bad block: freeze its coordinates and use Corollary III.109.7. Cross-block terms do not permit that immediate reduction. Natural first cases include interactions of lower degree in the bad block, bounded algebraic rank, or dependence through a few linear forms. The exact centre-count identity gives a baseline against which to measure how many progressions a perturbation destroys. A useful supersaturation theorem should record both the dimension of the added variables and the degree or rank of the interaction.

V.254.5 Closing a complete quantitative iteration

Question V.254.11 (A loss ledger from local theorems to a density increment). *Can the exact Fourier profile, the interpolation list improvement, and the applicable affine-recovery theorem be carried through one complete, explicitly parameterized Gowers density-increment step? What bound survives iteration?*

This is the necessary bridge to a stronger global theorem. The input and output at each lemma should record density, scale, codimension or partition complexity, exceptional mass, list size, and the precise characteristic assumptions. In particular, exponent-three rigidity cannot silently be inserted where the frequency target has a different torsion structure. Similarly, the Fourier neighbourhood theorem should not be used if the produced budgets lie outside its matrix-certificate region.

An attainable first result is a verified symbolic recurrence for a restricted class where all hypotheses match. One should then solve that recurrence and compare the resulting bound with its baseline, including lower-order dependence. This is more informative than multiplying apparent local gains whose hypotheses cannot all hold at once. Any claim of a new global bound must state the resulting function of N , k , and density and provide the full derivation.

V.255 Verification, reproducibility, and repository integration

Source 67, delivered Section 8.

V.255.1 Three levels of evidence

The package distinguishes proof, finite certificate, and diagnostic calculation.

1. **General proofs.** The Fourier moment argument, projection identities, autocorrelation inequality, interpolation upper and lower bounds, norm character estimate, and classification arguments are written out. Their general validity does not follow from the test suite, and no such inference is used.
2. **Finite certificates used in a theorem.** For the eleven small fields in the eight-term norm theorem, the displayed degree-nine polynomials are verified to be irreducible by exact polynomial arithmetic. The proof reduces irreducibility to $X^{p^9} \equiv X \pmod{F}$ and $\gcd(X^{p^3} - X, F) = 1$. The verifier also independently excludes factors of degrees one through four. This finite component is explicit and reproducible.
3. **Supplementary diagnostics.** Exhaustive maps test the graph identities and equality cases; finite interpolation examples test counts and uncovered-mass inequalities; exhaustive norm tables test normalizations. Floating-point Fourier checks compare the analytic answer with independently discretized linear programs. These calculations can catch mistakes but do not certify the infinite parameter families.

The main exact rigidity suite checks 115,623 maps, all 17,576 nonzero triples in $\mathbb{F}_3^3$ for the projection table, and 198,015 integer instances of the pointwise inequality. The additional order-thirteen suite enumerates all $3^{12} = 531,441$ ternary maps normalized by $f(0) = 0$; it confirms the maximum nonconstant energy 1645 and the 26 normalized one-point extremizers. This last enumeration is a check of an elementary analytic proof, not its substitute.

The interpolation suite includes all 28,561 bidegree-at-most- $(1, 1)$ polynomials over $\mathbb{F}_{13}$ in a small sharp example, deterministic generic-list constructions, greedy-order comparisons, and seeded multi-candidate cases. The norm suite verifies the eleven irreducibility certificates and eleven exhaustive norm-count cases. The Fourier suite tests 120 strict-branch examples with $R \leq 10$, an exact symbolic identity for the rational example, five independent grid linear programs, and nine finite-grid bounds. The numerical outputs are labelled as such in their JSON reports.

V.255.2 Files and reproducible commands

The distribution contains the complete single-file manuscript `gowers_exact_local_profiles.tex`, its compiled PDF, two vector figures, the section sources used for assembly, verification programs, their recorded outputs, and a source ledger. The main TeX file has no external bibliography dependency; its only external typesetting inputs are the two supplied PDF figures. The section sources are included for convenient selective integration.

To compile after extracting the archive, run

```
latexmk -pdf -interaction=nonstopmode -halt-on-error \
  gowers_exact_local_profiles.tex
```

The included `build.sh` performs the same operation. If `latexmk` is unavailable, running `pdflatex` three times resolves the contents and references. The verification commands, run from the extracted directory, are

```
python3 verification/verify_fourier.py
python3 verification/verify_rigidity.py
python3 code/verify_interpolation.py
python3 code/verify_order13.py
python3 code/norm_blocks_verify.py --full
```

Python’s standard library is sufficient for the exact rigidity, interpolation, and norm programs. The order-thirteen enumeration uses NumPy with integer arrays. The Fourier diagnostics use NumPy, SciPy, and SymPy. The optional figure regeneration script also uses Matplotlib. Run the exact programs without Python’s optimization flag: some use assertions as internal checks. The full norm and order-thirteen runs are more substantial than the remaining diagnostics. Their recorded results are included so that a reader can first inspect the certificates and then decide which checks to rerun.

V.255.3 Suggested integration points

Contribution	Existing location or dependency	Label namespace
Unequal Fourier budgets	Missing-interval discussion after Gowers’ Lemma 5.2; source 37, especially its unequal-budget question	fourier:
Ternary cubic rigidity and recovery	Graph-energy/affine-rigidity discussion; source 11 and the recent character/affine manuscript	cthree:
Exact candidate count	Shared-anchor step at Lemma 16.10; source 39’s candidate-count and further-research sections	interp:
Norm-block classification	Polynomial symmetric-colouring discussion; sources 10,14,22,38 and the Shi–Dong construction	norm:

The source ledger records the merged prior article’s inspected Git blob and hashes of the more recent working manuscripts. This is preferable to referring to an unqualified “latest” state during integration. The article is supplied for review and incorporation; no repository commit or external publication is represented as having been made by delivering this archive.

V.255.4 Formalization priorities

The most economical initial formalization target is the interpolation upper theorem. Its proof needs finite sets, a field, a linear space of base functions, and Lagrange interpolation. The greedy collision lemma and finite lower construction can follow, using elementary root bounds. These would give useful verified components without first formalizing Gowers’ entire proof.

The next target is the characteristic-three graph expansion and the two-color inequality. The latter reduces to a short real inequality under explicit integer constraints; the seven-case projection table can be encoded as finite linear algebra over $\mathbb{F}_3$. The global recovery theorem then needs finite Fourier analysis and an explicit treatment of characters of the wrong order. The order-thirteen cutoff is a compact additional trigonometric lemma.

For the norm section, separate the algebraic reduction and irreducibility certificates from the analytic character bound. The certificates can be checked by a small trusted polynomial arithmetic kernel. The general estimate may first be formalized conditional on the standard multiplicative Weil bound, with that dependency stated in the theorem. Finally, the Fourier section requires a finite supported-moment theorem or a formalization of its polynomial factorization proof. The finite matrix algebra, uniqueness, and variation formulas can be developed independently of that analytic input.

These are proposed formalization routes. The distribution contains mathematical proofs and executable checks, not Lean theorem files or claims of machine-checked proof.

V.256 Consequences, interfaces, and further research

Source 68, delivered Section 8.

V.256.1 Two direct interfaces with the Gowers argument

The energy $E(A)$ counts two-dimensional additive cubes in A . Thus Theorem II.297.1 gives, for every observed overlap,

$$\#\{\text{two-dimensional cubes in } A\} \leq |A|^3 - |A|r_A(d)(|A| - r_A(d)).$$

The multilevel version incorporates an entire overlap spectrum. These are structural bounds for a specified set; they can be combined with the uniformity-dependent cube bounds surrounding Gowers’ Lemma 3.10, but they do not change that lemma’s higher-dimensional uniformity parameter by themselves [1].

The restriction theorem has a direct moment interface. We include the argument to fix the density normalization and credit it as an existing ingredient [277, 279, 281].

Proposition V.256.1 (Existing moment lifting with the new filter). *Let B be a subset of a finite cyclic group, $b = |B| > 0$, and $\phi : B \rightarrow F$ a map into an abelian group. If at least γb^3 ordered*

additive quadruples are respected by ϕ , then $G_{16}(B, \phi) \geq \gamma^7 b^{15}$. Consequently, if $0 < \gamma, \eta \leq 1$ and

$$b \geq \left(\frac{4}{\gamma^7 \eta} \right)^{53},$$

there exists $B' \subset B$ with

$$T_{16}(B') \geq \left(\frac{\gamma^7 \eta}{4} \right)^{53} b^{15}, \quad G_{16}(B', \phi) \geq (1 - \eta) T_{16}(B').$$

Proof. First suppose the target is finite, and let Γ be the graph of ϕ in the finite product group of order M . Use the unnormalized Fourier transform $\widehat{\mathbf{1}}_\Gamma(\chi) = \sum_{x \in \Gamma} \overline{\chi(x)}$. Parseval gives $\sum_\chi |\widehat{\mathbf{1}}_\Gamma(\chi)|^2 = Mb$. Thus the weights $p_\chi = |\widehat{\mathbf{1}}_\Gamma(\chi)|^2 / (Mb)$ form a probability distribution. For the balanced $2s$ -tuple count,

$$G_{2s} = \frac{1}{M} \sum_\chi |\widehat{\mathbf{1}}_\Gamma(\chi)|^{2s} = b \sum_\chi p_\chi (|\widehat{\mathbf{1}}_\Gamma(\chi)|^2)^{s-1}.$$

For integer $s \geq 2$, convexity implies $G_{2s} \geq b(G_4/b)^{s-1}$. Taking $s = 8$ gives the first assertion. For an arbitrary target, the finite-quotient argument of Lemma II.311.7 preserves all relevant quadruple and sixteen-tuple discrepancies, so the same conclusion follows. Finally apply Theorem II.312.2 with $\alpha = \gamma^7$. $\square$

In Gowers' notation $b = \beta N$, so the size condition is explicitly $\beta N \geq (4/(\gamma^7 \eta))^{53}$. The difference between normalizing respected quadruples by b^3 and by N^3 must be maintained when this corollary is inserted into another statement. In particular, a density factor should not be removed by silently changing the meaning of γ .

V.256.2 A practical formalization order

The finite combinatorial part has a short route to a formal development. It is useful to keep all main inequalities over nonnegative integers or rationals until the last quadratic inversion. An implementation order is:

1. Define finite-support counting convolution and prove the union–intersection identity. Record (II.297.3) as an equality, then obtain the inequality by termwise positivity.
2. Prove the overlap triangle inequality and the integer form of the multilevel bound. For subgroup closure only two internal nodes are needed; its polynomial inequality can be proved without integration by clearing denominators in the chord calculation.
3. Construct H by the strict threshold $2r_A(d) > m$, prove it is a subgroup, and derive the occupancy identity over its cosets. Separate canonical-coset uniqueness from closest-coset uniqueness.
4. Formalize the complement expansion and the polynomial remainder. Prove the nested optimization either by the displayed derivatives or by an equivalent rational inequality after setting $s = \tau(s)(1 - \tau(s))$.
5. Formalize the finite degree-twelve coefficient convolution and integer certificates. Keep this independent of the analytic expectation theorem, so the finite arithmetic can be checked early.

6. Link the general finite-group expectation theorem to the existing ordered-tuple counts, and only then substitute the exponent-53 certificates.

The existing catalogue records exact companions for `lemma_3_10`, `lemma_9_3`, and `corollary_9_4`. This is a report about the inspected catalogue, not a new compilation claim. The tuple interface is the repository's `additiveTupleCount`, `phiAdditiveTupleCount`, and `GammaHomOfOrder` with order parameter eight. The proposed strengthenings should receive separate theorem names; the original statements and their current proofs need not be overwritten. No uncompiled Lean file is included as though it were a verified result.

V.256.3 Research questions left after the present results

The following questions are intended as specific next projects. They distinguish unknown sharp constants, extensions of the proved scope, and propagation tasks. The coefficient-one question, the quarter-defect coset existence threshold, the $2/9$ uniqueness threshold, the existence of a filter with cost below 52, and the leading growing-order phase diagram are answered above and are not listed again as open. The coordinate-splitting and Cartesian-structure parts of the compound ledger item `gsr:fz:q:II:cosetrangle` remain outside these answers.

Question V.256.2 (The discrete coset-distance envelope). *For a specified finite abelian group G , cardinality m , and defect budget $u \leq 1/4$, determine*

$$\Phi_{G,m}(u) = \max_{\substack{A \subset G, |A|=m \\ E(A) \geq (1-u)m^3}} \delta(A)$$

whenever the feasible class is nonempty. How do subgroup indices and the integer energy spectrum improve $\tau(u)$ between its attained punctured-coset parameters?

The equality classification supplies the exact candidates on the quadratic curve, while (II.298.5) records all losses away from them. A useful first case is a cyclic prime-power group with two or three possible nontrivial subgroup sizes. A complete answer should identify extremizers as well as the numerical envelope, and should distinguish the canonical distance δ_C from the true minimum $\delta(A)$ above the uniqueness threshold.

Question V.256.3 (Stability across the quarter-defect transition). *Fix $\kappa > 0$. Suppose $\kappa \leq q(d) \leq 1 - \kappa$ and $\varepsilon - q(d)(1 - q(d)) \leq v$. Is A within $O_\kappa(v^c)m$ edits, for an absolute $c > 0$, of one of the equality families in Theorem II.297.7? What are the optimal exponent and the necessary scale conditions?*

This asks for stability of the universal inequality itself, including the transition at $\varepsilon = 1/4$ where two-coset examples appear. The oriented puncture theorem answers a localized version with $\varepsilon < 1/4$ and a small gap relative to δ^3 . The exact remainder (II.297.4) suggests first proving that the mixed convolution v is nearly $\{0, k\}$ -valued and that w is nearly $\{0, t\}$ -valued, then converting those analytic conclusions into periodicity. The lower bound on $q(d)$ and $1 - q(d)$ prevents a vanishing overlap from making the hypothesis uninformative.

Question V.256.4 (Equality in the multilevel spectrum inequality). *Classify sets attaining (II.297.6) for their full distinct overlap spectrum. Which finite sequences of multiplicities can simultaneously make all the truncation-node inequalities sharp?*

For one internal node the answer is Theorem II.297.7; for a consecutive spectrum, integer intervals attain the bound. With several levels, equality forces the sum-multiplicity truncation curve to be exactly the prescribed piecewise linear interpolant. That is a strong compatibility condition between sum and difference representations, not merely a collection of scalar equalities. A natural starting class is a short arithmetic progression of cosets with one or more punctured fibers.

Question V.256.5 (A weighted coset theorem). *For finitely supported $0 \leq f \leq 1$ with $E(f) \geq (1 - \varepsilon)(\sum f)^3$, determine the sharp bound on*

$$\min_C \frac{\|f - \mathbf{1}_C\|_1}{\sum f}$$

in terms of ε . Does the same quarter-defect structural threshold have a sharp weighted distance formulation?

Proposition II.297.8 already gives a subgroup of large *minimum* overlaps. The missing step is a suitable replacement for the occupancy identity used for indicator sets. Simple tests should include a constant weight on a subgroup, a partially filled missing coset, and two fibers of unequal height. Any proposed formulation must retain the cap on f and the correct normalization under multiplication of all weights.

Question V.256.6 (Nonabelian analogues and asymmetric equality). *What sharp overlap-energy inequalities survive in finite nonabelian groups when left and right translations are distinguished? Within the abelian setting, classify all equality cases of the unequal-size mixed remainder (II.297.3).*

The abelian proof uses commutativity at the exact point where $\mathbf{1}_P * \mathbf{1}_C$ is replaced by $\mathbf{1}_C * \mathbf{1}_P$. This identifies the obstruction to a verbatim nonabelian extension. A useful initial problem imposes normality of the relevant stabilizer or studies unions of its cosets. For unequal sets, equality still forces the two nonnegative remainders to vanish, but full representation fibers no longer give the same reflection symmetry. A classification would strengthen the two-set form beyond its presently sharp universal coefficient.

Question V.256.7 (The exact sixteen-variable filter optimum). *Determine $\mathcal{I}_{16,2}$ or obtain a rigorous narrow interval for it. Can the degree-twelve construction be improved sufficiently to lower the clean common exponent below 53 after the size and collision constants are included?*

The exact upper certificate here is $\mathcal{I}_{16,2} < 103/2$; its actual filter cost is approximately 51.44505. This is an upper bound, not a global optimum. An effective search can parameterize nonnegative trigonometric polynomials as squared real polynomials or sums of squares, preserve nonnegative Fourier coefficients, and optimize gain and moments jointly. Rationalization should be followed by exact positivity, coefficient, and integer-moment certificates. A numerical grid check of $0 \leq K \leq 1$ is insufficient.

Question V.256.8 (The second-order principal cost). *For fixed a , does*

$$\mathcal{I}_{m,a} = \frac{m \log m}{2 \log a} + O_a(m)$$

hold? If not, identify the first unbounded correction. For $a = m^{c+o(1)}$, determine the first correction to $m(1 + 1/(2c))$ under a sufficiently precise specification of a .

The lower bound proved here has only an $O_a(m)$ loss. The upper construction has the remaining $O_a(m \log \log \log m)$ term. Thus the leading growing-order regime is resolved, but the triple logarithm is not known to be necessary. An answer should identify near-optimal filter shapes as well as scalar asymptotics. The relevant competing requirements are concentration of the coefficient distribution and near-uniformity of its m th powers among residue classes.

Question V.256.9 (Degree, principal cost, and finite-size cost together). *For fixed m, a , determine the tradeoff among filter degree D , principal cost $-\log P/\log(P/Q)$, and error cost $\log(R_2/P)/\log(P/Q)$. How large must D be to approximate the uniform theta-filter upper bound to a prescribed accuracy?*

The Fejér approximation proves an infimum statement without optimizing its degree. In an actual restriction theorem, increasing D increases the domain torsion allowance and the finite-size constant. Optimizing only P/Q can therefore give a worse usable theorem. The certified degree-twelve example is a concrete feasible point for this three-way optimization; a successful refinement would output all the constants required by Theorem II.311.6.

Question V.256.10 (Using the observed discrepancy-order distribution). *Suppose the bad tuples have known counts u_a by finite discrepancy order, together with a count for infinite order. Optimize a restriction certificate using the actual principal sum $\sum_a u_a Q_{m,a}^k$ rather than the order-two worst case. How much can independent heterogeneous filters improve the complete finite-size guarantee for such a profile?*

The uniform theorem gives the cost for each single order, but it does not solve a multi-order optimization. A filter adapted to one residue class can be inefficient for another. The natural optimization must retain the weighted nonprincipal and collision errors as well. Good benchmark inputs would separate a small population of involuntary defects from a large population of high-order defects.

Question V.256.11 (Methods beyond the principal positive-filter barrier). *Can one obtain retention exponents of order $O(m)$ for a single balanced relation on every cyclic group with arbitrary abelian targets, with useful explicit finite-size thresholds, by a method whose guarantee is not controlled solely by the principal positive Fourier spectrum?*

For fixed possible discrepancy order, the proved principal cost grows like $m \log m$, including for independent heterogeneous layers. This leaves adaptive filters, correlated filters, signed Fourier coefficients with controlled cancellation, and prior structural extraction as distinct possibilities. The source-59 character-separation results are already available [282]; a new project should combine them with the present tuple-retention objective, not repackaging character separation itself as an unanswered question.

Question V.256.12 (Several relations with decaying collision errors). *Extend the source-58 expectation estimate to a fixed integer relation matrix, selecting one common subset and counting every ordered tuple. Can the nonprincipal error retain decay in the number of layers while its domain dependence is expressed through Smith invariants or coefficient-minor kernels?*

The next tractable case is two independent balanced relations. One must track the rank gained by each added Fourier row and preserve the geometry of tuples in B^m , including identifications of coordinates. General system-level Fourier restriction and the repository's finite-field rank methods are predecessors, so the specific goal is the simultaneous decaying-error, composite-modulus,

common-subset statement. A hybrid primary-component argument may also connect cyclic bounds with useful bounds on large groups of bounded exponent.

Question V.256.13 (Construction and global propagation). *Can a subset meeting the explicit certificate be found and verified without enumerating all b^m tuples? Separately, which later steps of the Gowers induction dominate after substituting the new local restriction and stability constants?*

For construction, a fixed finite Fourier degree allows the phase integrals to be replaced by a sufficiently fine exact root-of-unity average, and conditional expectation gives a finite exhaustive procedure. Efficient convolution, partial tuple statistics, and certified algebraic arithmetic are possible routes to a useful algorithm; no such complexity bound is proved here. For global propagation, the appropriate deliverable is a dependency ledger carrying densities, quality parameters, domain sizes, and normalizations through every use of the restriction statement and the higher arrangement lemmas. A local exponent improvement does not determine the final Szemerédi bound until those later inequalities are recomputed.

V.256.4 Suggested order of attack

The most immediate projects are exact finite filter optimization, discrete coset envelopes in small subgroup lattices, and formalization of the nonnegative energy remainder. They have explicit certificates and finite benchmarks. The next layer is stability at the quarter-defect transition and the second-order principal cost; both have a clearly identified gap between the available upper and lower mechanisms. Several-relation restriction and a revised global parameter chain are larger projects whose hypotheses should be fixed before numerical optimization begins.

V.257 Verification, provenance, and integration notes

Source 68, delivered Section 9.

V.257.1 What has been proved and what has been checked

The theorems for arbitrary finite sets, arbitrary abelian ambient groups, and unbounded parameters are justified by the written proofs. Finite computations serve a different purpose: they audit identities, small torsion, endpoint conventions, equality classifications, and explicit arithmetic constants. They are not substitutes for those proofs.

The preparation used parallel mathematical derivations and internal cross-checks. In particular, the pointwise gap, subgroup threshold, coset rounding identity, puncture derivatives, and rational filter certificates were reviewed independently within the same AI-assisted preparation process. This does not constitute external referee review. No Lean executable was available in the preparation environment, and no theorem in this supplement is represented as newly kernel-checked.

The following completed runs are included as machine-readable JSON reports. Counts refer to the enumerated assertions in their own suites; they should not be added as if all were distinct input sets.

Exact check	Cases
Nonempty sets in the main energy suite	394,229
Shifted pointwise inequalities and exact remainders	6,168,924
Full-spectrum multilevel inequalities	394,229
High-energy subgroup and canonical-coset cases	20,694
Pointwise equality cases (including zero defect)	38,713
Canonical-envelope equality cases	5,672
Ordered mixed-set pairs	738,884
Translated mixed-set cases	6,404,265
Rational-weight functions	7,726
Shifted weighted exact remainders	58,794
Weighted high-energy subgroup cases	1,015
Ordered mixed-weight pairs	13,544
Sets and models in the separate puncture suite	65,572
Coset-remainder checks in that suite	32,960
Nested-profile instances, including inner endpoint	778
Oriented-profile instances, including inner endpoint	11,420
Degree-twelve exact integer certificates	17

The energy suite enumerates every nonempty subset of $C_n = \mathbb{Z}/n\mathbb{Z}$, $2 \leq n \leq 16$, and of

$$C_2^2, C_2^3, C_3^2, C_4 \times C_2, C_4^4, C_4^2, C_8 \times C_2, C_4 \times C_2^2.$$

All ordered mixed pairs are included for those groups of order at most nine. The rational-weight suite uses values $\{0, 1/2, 1\}$ on $C_1, \dots, C_6, C_2^2, C_2^3$, excluding the zero function, and all ordered pairs on included groups of order at most four. Integer numerators avoid floating-point acceptance tests.

The separate puncture suite enumerates all nonempty subsets of C_n for $1 \leq n \leq 15$, then adds exact nested constructions, strict inequality cases, actual outliers, nontrivial stabilizers, and proper outer cosets, including examples in C_{2000} . Its 291 strict-inner nested instances include 63 beyond the outer $2/9$ uniqueness range. It also checks 487 inner-endpoint nested instances. The oriented totals are 6,754 strict-inner and 4,666 inner-endpoint instances. The examples in C_4 , C_6 , C_8 , and C_{12} separately test closest-coset nonuniqueness, failure at the outer quarter endpoint, attainment at the inner quarter endpoint, and nonunique optimal holes in the enlarged range.

For the filter, all acceptance conditions are comparisons of explicit integers, including positivity, symmetry, monotonicity, normalization, and the seven moment inequalities printed in Proposition II.312.1. The floating-point values in (II.312.6) are descriptive only. The filter verifier also succeeds with Python optimization enabled, because its mathematical acceptance checks do not rely on `assert`. The other exhaustive verifiers should be run with ordinary Python, with assertions enabled.

V.257.2 Reproducing the checks and the article

The package contains the full standalone `gowers_sharp_energy_restriction.tex`, its compiled PDF, the modular source `main.tex` with `sections/`, the figure and its generator, verification programs, and their reports. The standalone source has its section inputs expanded; the optional figure remains a separate included PDF in `figs/`. Keep that folder beside the source when rebuilding the supplied layout. The figure can be omitted without affecting any proof.

From the extracted package directory, the verification commands are:

```
python code/verify_energy.py
python code/verify_energy.py --weights-only
python code/verify_punctures.py > data/puncture_checks.json
python code/verify_filter.py > data/filter_checks.json
```

The programs use only the Python standard library. The initial full energy run took approximately 90 seconds in the preparation environment; timings depend on the machine. Rebuilding the article requires an ordinary pdfLaTeX installation and the packages named in the preamble. The included Makefile runs pdfLaTeX three times to stabilize a clean build. Matplotlib and NumPy are required only to regenerate the already supplied figure.

The delivered PDF was compiled from the supplied source, checked for unresolved references and layout warnings, rendered, and visually inspected. The archive is accompanied by file checksums and integration instructions. A build or a finite verifier succeeding establishes the stated computational check; it does not upgrade the mathematical status to a formal proof.

V.257.3 Pinned source ledger

All current-tree paths below are relative to `Combinatorics/Ramsey/` in [the pinned ProveIt Ramsey tree](#).

1. `Papers/sz-thm-gowers-proof/sz-thm-gowers-proof.tex`. The repository transcription was checked against the primary published paper for the statements and normalizations used here. The relevant interfaces are the additive-energy discussion around Section 7, Lemmas 9.1–9.3, Corollary 9.4, and the two-dimensional cube interpretation near Lemma 3.10.
2. `Research/GowersSzemerédi/local-quantitative-refinements/article.tex` and its `README.md`. The consolidated manuscript covers sources 01–39. We inspected the relevant energy and restriction passages, especially sources 04, 14, and 17, and the consolidated research questions. This was a targeted mathematical audit, not a line-by-line reproof of all earlier material.
3. Separately placed records for sources 40, 45, 47, 52–59 were inspected for overlapping claims and source status. The full manuscripts used for the closest comparisons were source 45, *Rigidity gaps in Gowers’s Szemerédi framework*; source 58, *Polynomial restriction on every cyclic group*; and source 59, *Sharp Character Separation and Affine Rigidity in Gowers’ Inverse Steps*.
4. Source 45’s full archive was retrieved at commit `8ed496b14e46aad80a50fa7c34ea50d3c66a9323`, path `docs/incoming/Gowers_Rigidity_Gaps_Research.zip`. Sources 58 and 59 were retrieved at commit `f8bc5e2ec281fcee02c0fc5ef717a775ee7e4442`, paths `docs/incoming/gowers_torsion_uniform_restriction.zip` and `docs/incoming/gowers-character-affine-refinements.zip`. These historical archive pins identify the full inspected documents; their own older inspection pins do not replace this article’s current-tree comparison pin.
5. The formal-status catalogue was consulted for the existing exact companions and interface names. Its status is reported as found; no formal-status file is modified by this contribution.

Source 45 supplies the prior factor-two gap, small-range coset envelope, complement identity, and nested repair profile. Source 14 supplies the earlier sharp quadratic profile and the input-triple counting used in the stronger remainder. Source 58 supplies the decaying-error restriction theorem, all-cyclic torsion handling, exponent 54, and fixed-order principal cost theorem. Source 59 already supplies broad character separation and affine-rigidity results; those are not claimed as new here. The present claims of improvement are exactly those in the table of Section F.73.

V.257.4 Integration safeguards

Assign a new source identifier through the repository's normal intake process. Prefix this article's labels and bibliography keys when incorporating its sections into the consolidated report. Preserve the original source attributions and the explicit hypotheses: ordered tuples with repetitions; arbitrary abelian targets; cyclic domains for the uniform constant 144; the large point-count threshold; canonical versus globally closest cosets; strict versus closed quarter endpoints; and principal-spectrum versus complete subset-selection cost. The energy, puncture, finite filter, and asymptotic sections can be integrated independently according to their proved dependencies. The repository itself has not been changed by preparing this package.

V.258 Finite certificates and implementation interfaces

Source 69, delivered Section 10.

V.258.1 Coefficient-table algorithms

A symmetric degree- d tensor is represented by its $\binom{n+d-1}{d}$ coefficients on nondecreasing basis tuples. From this table the following operations are finite linear algebra over $\mathbb{F}_2$.

First form the matrix A_R whose columns correspond to a basis direction in the first frozen slot and whose rows correspond to a tuple of $d - 4$ remaining frozen indices and an alternating pair $a < b$. Its entries are

$$(A_R)_{(\mathbf{w},a,b),z} = T(a, a, b, z, \mathbf{w}) + T(a, b, b, z, \mathbf{w}). \quad (\text{V.258.1})$$

Its kernel is $R(T)$ and its rank is $\rho(T)$. There are $\binom{n}{2} \binom{n+d-5}{d-4}$ rows before redundant rows are removed. Next form the matrix A_P with a row for each frozen tuple of length $d - 3$ and each basis vector a , and a column for the other alternating variable b :

$$(A_P)_{(z,a),b} = T(a, a, b, z) + T(a, b, b, z). \quad (\text{V.258.2})$$

It has $n \binom{n+d-4}{d-3}$ rows. The stacked matrix $\begin{pmatrix} A_R \\ A_P \end{pmatrix}$ has kernel $K(T)$ and rank $m_{\text{eff}}(T)$.

Gaussian elimination yields both radicals. A complement of $K(T)$ gives a linear quotient q and section j . Restricting the tensor to that complement constructs $T_0 = j^*T$. To certify the result, verify $qj = I$ and verify on basis tuples that $\Phi_T = q^*\Phi_{T_0}$. The difference is integrable, and its support coefficients construct a primitive by (IV.59.7). All equality claims in this certificate are exact bit identities.

With naive elimination, the matrix work uses $O((r_R + r_P)n^2)$ field operations for the displayed numbers of rows, in addition to coefficient indexing and output construction. This is a bound for an

explicit coefficient-table input. It is not a polynomial-time claim for tensors supplied by a succinct circuit while the degree varies.

V.258.2 Certifying the extremal flag

When $\rho(T) = 1$, row reduction identifies the unique nonzero u with $R(T) = \ker u$. Choose w with $u(w) = 1$ and evaluate all frozen slots at w to obtain Ω . Define $v(x) = \Omega(x, w)$. The Bianchi identity gives $\Omega = u \wedge v$, and v is independent of u because the defect is nonzero. Verify that $T - C_d^{u,v}$ has support-dependent coefficients and construct its primitive. This certificate identifies an exact extremal gauge class without searching over complex-valued functions.

V.258.3 What the delivered code does

The program `code/verify.py` implements exact coefficient arithmetic, contractions, integrability tests, the two radical matrices, the sharp model construction, and direct selected/cube energy calculations in small groups. It is a reproducibility verifier, not a production inverse-theorem package. Its finite scope and output counts are recorded in Section V.260 and `data/verification.json`. No value of a finite optimization is silently promoted to a general theorem.

V.259 Relation to Gowers’s argument and the ProveIt formal project

Source 69, delivered Section 11.

V.259.1 The mathematical interface

The connection to Gowers’s Section 17 is the same normalized object: an average of squared Fourier coefficients of iterated derivatives. Structured frequencies are used to remove a polynomial phase. Our conclusion is an exact identity after a *certified integrable* symmetric Boolean tensor has been supplied. It addresses the small-characteristic obstruction that ordinary symmetry alone does not remove.

There are important differences from the existing cyclic statement. The inspected `proposition_17_2` in the repository [285] in `Sections17_18.lean` assumes prime modulus, factorial invertibility, and a multiaffine frequency map in the repository’s sense. It records a phase of degree at most $k + 1$, consistent with applying the preceding finite-difference lemma to $x_{k+1}\sigma(x_1, \dots, x_k)$. The original printed statement says k ; we use the corrected formal interface rather than that typographical degree.

Here the ambient group is $\mathbb{F}_2^n$, the tensor is homogeneous multilinear and symmetric in all d slots, and the primitive takes values in $\mathbb{R}/\mathbb{Z}$. Setting $d = k + 1$ aligns the number of derivative directions and the norm power. It does not turn this result into a literal replacement for the cyclic theorem, nor justify importing its multiaffine hypotheses unchanged. The theorem is a related sharp small-characteristic phase-integration result.

V.259.2 Formalization milestones

A suitable new namespace should first separate symmetric multilinear maps from the repository’s multiaffine predicate. The following dependency chain is sufficient:

Component	Required mathematical content
Finite calculus	Cube expansion, commutation, exact slicing, and gauge invariance, all allowing zero amplitudes.
Boolean integration	Repetition identity, support-indexed coefficients, explicit circle-valued primitive, and the dimension of $\mathcal{I}_d$.
Cubic base	Finite-Hilbert-space projective orthogonality, shear identity, and the alternating-rank bound.
Canonical branch	Derivative periods, constrained pure-power cap, rank-one classification, and canonical recurrence.
Universal induction	Kernel counting, nonintegrable contractions, and $3d > 2d + 2$ for $d \geq 4$.
Model theorem	Complete radical, quotient factorization, integrable diagonal lowering, rank-nullity, and support-basis sharpness construction.

The main all-degree induction is short once these interfaces are available. The model theorem is largely finite linear algebra and can be formalized separately from the spectral analysis. A verified implementation should reconstruct all finite certificates from definitions, not treat the delivered JSON as a trusted theorem database. No Lean code is delivered, and no build or update of the repository’s proof-status ledger is claimed.

V.259.3 What does not follow

The near-one thresholds here do not themselves turn a small positive Gowers norm into a polynomial phase. There is no new low-energy symmetrization or localization argument, no control of Bohr-set boundary losses, and no iteration producing a new bound for progression-free subsets of an interval. For perspective, the global bounds of Leng–Sah–Sawhney [2] involve additional quantitative inverse-theorem machinery; the present estimates are not a claimed improvement to those bounds. Nor does the effective tensor quotient prove that every optimizer descends to it. The canonical all-dimensional equality is justified by its separate period argument, not by an unsupported quotient principle.

V.260 Exact verification and source provenance

Source 69, delivered Section 12.

V.260.1 The finite checks actually performed

The shipped verifier uses Python’s standard library, integer bit matrices, Gaussian integers, and exact rational arithmetic. Failures raise exceptions; checks are not disabled by `python-0`. The default run records:

1. The canonical recurrence and strict induction separation in every degree from 4 through 1000, with rational arithmetic.
2. Every symmetric binary tensor in degrees 3 through 12: $16 + 32 + \dots + 8192 = 16368$ tensors. Exact constant-function cube energies are compared with the stated threshold and the rank-sensitive bound. These are not optimizations over all functions.
3. All 32768 symmetric quartic tensors on $\mathbb{F}_2^3$. Both defect-flattening matrices are computed to check the effective dimension and rank-one classification. The counts are those printed in Section IV.63.
4. The sharp model construction for each $d = 4, 5, 6, 7$ and $r = 1, 2, 3$. The full coefficient tables, frozen radicals, complete radicals, and support-basis maps Ψ are checked. The largest example has ten coordinates, degree seven, and 11440 unordered coefficients.
5. 100 fixed-radical normal-form dimension checks, covering $n = 2, \dots, 6$, $d = 4, \dots, 7$, and every complement dimension $r = 0, \dots, n$. The constraints are row-reduced as equations on the full tensor coefficient space and compared with (IV.64.7).
6. 2040 exact tests of diagonal lowering and ordinary contractions of integrable tensors, in dimensions one through four and degrees three through seven, using a fixed seed for the support coefficients.
7. 46 exact complex-function tests of the selected/cube identity and slicing, in $(n, d) = (2, 3), (2, 4), (3, 3)$. Function values are Gaussian integers divided by two and include zero and nonunit amplitudes. The same run checks 46 canonical energy bounds, 112 constrained pure-power derivative bounds, and 5312 vector identities of each type: translation-conjugation and even-direction periodicity.

The cubic and arbitrary-dimensional analytic claims rest on their proofs, not on a finite phase grid. The coefficient censuses can detect sign and rank errors, but cannot certify a continuous extremum by themselves.

V.260.2 Primary and project sources

The original Gowers paper was inspected at the page image containing Proposition 17.2, printed page 578. Tidor’s Definition 3.1 and Proposition 3.5 were checked in the primary paper. The immediate project baseline [284] was read from the Library source titled *Exact Boolean Obstruction Energies*, together with its repository source audit. Its displayed main bounds, pure-defect classification, canonical proof, and explicit higher-degree question are the comparison points. Source 43 [283] is credited for its preceding cubic and shear framework, as identified by that audit. The formal interface was fetched from the observed repository revision

0d93c1eac2e20c313e7ed0cb7cb74ff705b00101.

The inspected `Sections17_18.lean` file has blob

18495f50475835ca1c497d85b0052215f865b8ca.

The preceding source 47 records its own comparison revision `5c9a442d263632fdcd4e9690154c12c2dcc70b53`; these are separate provenance records, not claims

that the entire repository was built at either revision. A targeted search did not establish independent literature priority for the new theorems. Accordingly, no claim of an exhaustive novelty search or external peer review is made.

V.261 Further research questions and topics

Source 69, delivered Section 13.

The following are proposed continuations. They are not assumptions in the proofs and are not presented as a list of previously published famous open problems.

Research question V.261.1 (The second extremal level). For fixed $d \geq 4$, determine

$$\sup_{n \geq 2} \max_{T \in \text{SM}_d(\mathbb{F}_2^n) \setminus \mathcal{I}_d(\mathbb{F}_2^n)} \mathcal{M}(T),$$

$\rho(T) \geq 2$

with the maximum over an empty stratum omitted. Is the upper bound $1 - 3d/2^{d+1}$ attained for any or all degrees? The deficit identity (IV.62.3) forces simultaneous extremality of almost every nonintegrable contraction in an equality example, while the derivative constraint may obstruct this compatibility.

Research question V.261.2 (Functions at the canonical endpoint). Classify all one-bounded f with $\mathcal{E}_{C_d^{u,v}}(f) = b_d$, modulo translations and multiplication by phases of degree at most $d - 1$. The present theorem classifies tensor classes, not functions. The nonnegative slice deficits, the fixed-derivative zero-frequency constraints, and equality in projective orthogonality provide concrete necessary conditions.

Research question V.261.3 (Dimension-free phase stability). Does $\mathcal{E}_{C_d^{u,v}}(f) \geq b_d - \delta$ force f to be close in normalized L^2 to the set of exact extremizers, with a modulus depending on d and δ but not on n ? Compactness only gives a modulus at a fixed ambient dimension. A proof would have to retain common-function constraints across many derivative slices.

Research question V.261.4 (Stabilization outside the canonical class). For which symmetric tensors T_0 and surjections $q : V \rightarrow W$ is $\mathcal{M}(q^*T_0) = \mathcal{M}(T_0)$? The canonical family has this property by the period argument. The sharp quotient theorem alone does not establish it for a noncanonical binary tensor or for the five-coordinate quartic (IV.65.5). A counterexample would identify a precise limit to tensor-model reduction as an optimization method.

Research question V.261.5 (Energy versus the sharp model size). Determine the best upper bound on $\mathcal{M}(T)$ jointly in terms of $\rho(T)$ and rank $\Psi = m_{\text{eff}}(T) - \rho(T)$. The sharp dimension examples saturate the algebraic bound, but no claim is made that they saturate an energy bound. Additional active alternating directions might force stronger spectral loss than (IV.62.2) sees.

Research question V.261.6 (Classification of maximal model tensors). Classify the tensors with $m_{\text{eff}}(T) = B_d(\rho(T))$ modulo integrable tensors and linear coordinate changes. The normal form in Theorem IV.64.4 reduces this to an action on its two parameters. The dimension proof shows that their radical map $\Psi : R(T)/K(T) \rightarrow \mathcal{I}_{d-2}(H)$ must be surjective. Which additional compatibilities distinguish the construction in Section IV.65 from other surjective cases?

Research question V.261.7 (Odd-characteristic counterparts). Over $\mathbb{F}_p$ the nonclassical repetition relation moves p occurrences rather than two. Find the appropriate contraction radical, a

compatible lowering operator, and the exact analogue of the nonintegrable energy threshold. The Boolean derivative-period proof does not transfer unchanged: translation by a direction no longer conjugates the derivative product.

Research question V.261.8 (Mixed functions and incomplete symmetry). What survives when the same function is not used at every cube vertex, or when the frequency tensor is not fully symmetric? In these settings the exact period identities and the contraction-integrability kernel may no longer suffice. A correct extension should explicitly retain the lost compatibility conditions, rather than apply the present threshold without its hypotheses.

Research question V.261.9 (Local versions and global bookkeeping). Can a robust local substitute for the contraction identity be established on Bohr sets, subprogressions, or high-density subsets, with an explicit boundary error? To influence a full density-increment proof, such a result must work at the energy scales actually produced by the inverse step, not only near one, and every localization and iteration loss must be propagated.

Research question V.261.10 (Verified finite certificates). Formalize the support-indexed primitive, diagonal lowering, and the two radical matrices first, then verify the exact minimal-quotient algorithm. Can the spectral cubic bound be connected to this algebraic layer with a small finite-dimensional Hilbert-space interface, so that the all-degree induction becomes a short kernel-checked theorem?

V.262 Conclusion

Source 69, delivered Section 14.

The all-degree extremal problem closes because nonintegrability of a contraction is an exact algebraic event with a subspace as its complement. When that subspace is a hyperplane, the obstruction is canonical; every other nonintegrable tensor has at least three quarters of its contractions outside the integrable class. The lower-degree energy loss then dominates the canonical loss strictly. This yields the exact threshold and the flag classification in every degree.

The same radical has a second role. Its interaction with the alternating pair produces an integrable-valued linear map, whose rank measures all additional active coordinates. The resulting dimension bound is exact and is attained by a support-basis construction. These two results sharpen the Boolean phase-integration part of the Gowers framework while keeping the remaining analytical and formalization boundaries explicit.

V.263 Reproduction and package layout

Source 69, delivered Appendix A.

The package contains this standalone `article.tex`, its compiled `article.pdf`, `code/verify.py`, the generated `data/verification.json`, and source, formalization, and build notes. From the package root run

```
python3 code/verify.py --max-degree 12
latexmk -pdf -interaction=nonstopmode -halt-on-error article.tex
```

The verifier requires Python 3.10 or later and no third-party Python package. The source uses standard pdfLaTeX packages, an embedded bibliography, no external figures, and no shell escape. The distributed PDF is compiled from the accompanying source. Previous manuscripts, third-party papers, and font files are not redistributed.

References

Merged from the five sources' bibliographies. A work cited by several sources appears once, with each source's annotations; the repository entries are kept per source, because each records that source's pin.

- [1] W. T. Gowers, *A new proof of Szemerédi's theorem*, Geometric and Functional Analysis **11** (2001), 465–588. doi:10.1007/s00039-001-0332-9 (02, 03, 04). Public PDF: <https://www.cs.umd.edu/~gasarch/TOPICS/vdw/sz-thm-gowers-proof.pdf> (01, 03, 04; source 01 gives the path `users/gasarch`). Source 01: relevant components are Corollaries 2.4–2.5, Lemma 17.4, Corollaries 17.5–17.6, Proposition 17.7, and the partition interface of Section 18; the public PDF was consulted. Source 05: the theorem numbering used for comparison is that of this paper; finite corrections in the repository are identified separately ([repo05*] below).
- [2] J. Leng, A. Sah and M. Sawhney, *Improved bounds for Szemerédi's theorem*, arXiv:2402.17995, version 2, 29 February 2024. <https://arxiv.org/abs/2402.17995> (all five sources). Source 05: the comparison uses Theorem 1.1 and the simultaneous-localization discussion in Sections 1–2; it does not import their quantitative theorem into any proof.
- [3] J. Leng, A. Sah and M. Sawhney, *Quasipolynomial bounds on the inverse theorem for the Gowers $U^{s+1}[N]$ -norm*, arXiv:2402.17994 (version 3, 10 August 2024, per source 02). <https://arxiv.org/abs/2402.17994> (02, 03, 04).
- [4] B. Green and T. Tao, *An inverse theorem for the Gowers $U^3(G)$ norm*, Proceedings of the Edinburgh Mathematical Society **51** (2008), 73–153. arXiv:math/0503014. The arXiv abstract includes a December 2023 correction concerning Proposition 3.2; no use of that proposition is made (source 02).
- [5] B. Green and T. Tao, *Yet another proof of Szemerédi's theorem*, arXiv:1002.2254 (2010). <https://arxiv.org/abs/1002.2254> (source 03).
- [6] B. Green and T. Tao, *New bounds for Szemerédi's theorem, III: A polylogarithmic bound for $r_4(N)$* , 2017. <https://arxiv.org/abs/1705.01703> (source 04).
- [7] T. Eisner and T. Tao, *Large values of the Gowers–Host–Kra seminorms*, arXiv:1012.3509 (source 02).
- [8] P. Baldi and R. Vershynin, *Polynomial threshold functions, hyperplane arrangements, and random tensors*, arXiv:1803.10868v3 (2019). Section 1.4 records the classical linear-threshold estimates; Theorem 1.1 supplies the polynomial-threshold asymptotic used here. <https://arxiv.org/abs/1803.10868v3> (source 01).
- [9] Yu. A. Zuev, *Asymptotics of the logarithm of the number of Boolean threshold functions*, Soviet Math. Dokl. **39** (1989), 512–513; Russian original in Dokl. Akad. Nauk SSSR **306** (1989), 528–530. Classical attribution and estimate consulted through [8] (source 01).
- [10] J. Kahn, J. Komlós and E. Szemerédi, *On the probability that a random ± 1 -matrix is singular*, J. Amer. Math. Soc. **8** (1995), 223–240. The Section 4 threshold-count estimate is used as recorded in [8] (source 01).
- [11] J. Fox and B. Sudakov, *Dependent random choice*, arXiv:0909.3271 (2009). <https://arxiv.org/abs/0909.3271> (source 03).
- [12] C. Reiher and T. Schoen, *Note on the theorem of Balog, Szemerédi, and Gowers*, Combinatorica **44** (2024), 691–698. doi:10.1007/s00493-024-00092-5; <https://arxiv.org/abs/2308.10245> (source 03).
- [13] M. Jaber, Y. P. Liu, S. Lovett, A. Ostuni and M. Sawhney, *Quasipolynomial bounds for the corners theorem*, arXiv:2504.07006 (2025), version 2. <https://arxiv.org/abs/2504.07006> (source 03).
- [14] F. Manners, *True complexity and iterated Cauchy–Schwarz*, 2021. <https://arxiv.org/abs/2109.05731> (source 04).

- [15] [source 01] V. Reshetnikov, *ProveIt: Combinatorics/Ramsey*, repository snapshot 3a25020a654d196f7f1aa2e37649752c000d0e65, inspected 6 October 2026. <https://github.com/VladimirReshetnikov/ProveIt/tree/3a25020a654d196f7f1aa2e37649752c000d0e65/Combinatorics/Ramsey>.
- [16] [source 01] *Definitions.lean*, in the snapshot [15], under *Combinatorics/Ramsey/Lean/GowersSzemerédi/*. Inspected definitions: partition, density, balanced function, unnormalized Fourier transform, natural progression, and average cell size.
- [17] [source 01] *Proofs02DensityIncrement.lean*, in the snapshot [15]. Inspected header and preparatory density and positive-mass lemmas.
- [18] [source 01] *Proofs17FinitePatterns.lean*, in the snapshot [15]. Inspected header and finite-dimensional sign-family infrastructure.
- [19] [source 01] *Sections17_18.lean*, in the snapshot [15]. Inspected floor-pattern specifications, affine patterns, partition uniformity, and the editorial dependency warning attached to Theorem 18.1.
- [20] [source 02] V. Reshetnikov, *ProveIt repository*, inspected October 6, 2026. Paper transcription at <https://github.com/VladimirReshetnikov/ProveIt/tree/main/Combinatorics/Ramsey/Papers/sz-thm-gowers-proof>; Lean development at <https://github.com/VladimirReshetnikov/ProveIt/tree/main/Combinatorics/Ramsey/Lean/GowersSzemerédi>. Specific inspected blobs are listed in source 02’s Appendix A (Section V.11). These links follow the moving branch *main*; the blob identifiers are the permanent record.
- [21] [source 03] V. Reshetnikov, *ProveIt: Gowers paper transcription and Lean development*, repository revision 3a25020a654d196f7f1aa2e37649752c000d0e65, inspected 6 October 2026. [Pinned paper directory](#); [pinned Lean directory](#).
- [22] [source 03] V. Reshetnikov, *Gowers proof status inventory*, *Combinatorics/Ramsey/gowers-proof-status.json*, at the revision in [21]. [Pinned source inventory](#).
- [23] [source 04] Vladimir Reshetnikov and contributors, *ProveIt: paper transcription and Lean development for Gowers’s proof of Szemerédi’s theorem*, GitHub repository. Inspected 6 October 2026, snapshot 327773149bd272239f1f6c0015f76d5097512c88. <https://github.com/VladimirReshetnikov/ProveIt/tree/327773149bd272239f1f6c0015f76d5097512c88/Combinatorics/Ramsey/Papers/sz-thm-gowers-proof>. <https://github.com/VladimirReshetnikov/ProveIt/tree/327773149bd272239f1f6c0015f76d5097512c88/Combinatorics/Ramsey/Lean/GowersSzemerédi>.
- [24] [source 05] V. Reshetnikov, *ProveIt repository*, edited transcription *sz-thm-gowers-proof.tex*, in *Combinatorics/Ramsey/Papers/sz-thm-gowers-proof*. Inspected 6 October 2026; returned blob identifier 41085a0efde8c1932e86e80c791984f52841c22e. [Repository paper directory](#).
- [25] [source 05] V. Reshetnikov, *ProveIt repository*, *Proofs02Partition.lean*, in *Combinatorics/Ramsey/Lean/GowersSzemerédi*. Returned blob identifier 291b1494ab5af839456dcac35f4bb447d4c90844. [Commit-pinned partition module](#).
- [26] [source 05] V. Reshetnikov, *ProveIt repository*, *Proofs02DensityIncrement.lean*, in the same directory. Inspected portions include the balanced-function identities and the large-scale proof of Corollary 2.5. Returned blob identifier 85b5c3ad8769b9769e1939ae5c05f0757f0678b5. [Density-increment module](#).

Added in the second write (sources 06–10, 6 October 2026). Works above that sources 06–10 also cite: [1] (06, 07, 08, 09, 10; source 07 adds “Section 10, particularly Lemmas 10.6, 10.8, and 10.9”, source 08 “in particular, Lemmas 3.7–3.10” and the visual check of printed pages 483 and 485); [2] (all five); [3] (06); [9] (06, with the MathNet record <https://www.mathnet.ru/eng/dan7111>); [8] (06, as the journal version: SIAM J. Math. Data Sci. **1** (2019), 699–729, [doi:10.1137/19M1257792](https://doi.org/10.1137/19M1257792)); [12] (10, Theorem 1.2); [6] (10); [14] (08). The entries below are new.

- [27] M. Ben-Or, D. Coppersmith, M. Luby and R. Rubinfeld, *Non-abelian homomorphism testing, and distributions close to their self-convolutions*, Random Structures & Algorithms **32**(1) (2008), 49–70; earlier version ECCC TR04-052 (2004). [Author’s bibliographic record](#). For the relevant structure theorem and majority-difference description see also [28], Theorem 4.2 (source 07).
- [28] P. Diaconis, X. Shao and K. Soundararajan, *Carries, group theory, and additive combinatorics*, arXiv:1309.0434 (2013), version 1. <https://arxiv.org/abs/1309.0434v1> (source 07).
- [29] W. T. Gowers and J. Wolf, *The true complexity of a system of linear equations*, Proceedings of the London Mathematical Society **100** (2010), 155–176. [arXiv:0711.0185](https://arxiv.org/abs/0711.0185) (source 08).
- [30] S. Cho, D. Conlon, J. Lee, J. Skokan and L. Versteegen, *On norming systems of linear equations*, 2024. [arXiv:2411.18389](https://arxiv.org/abs/2411.18389) (source 08).
- [31] H. Hatami, *On generalizations of Gowers norms and their geometry*, arXiv:0903.3237 (2009). <https://arxiv.org/abs/0903.3237> (source 09).
- [32] L. Milićević, *A quasipolynomial inverse theorem for the $U^k(\mathbb{F}_p^n)$ norm in the high characteristic*, arXiv:2609.33733 (2026). <https://arxiv.org/abs/2609.33733> (source 09; existence checked at intake).
- [33] R. Shi and Y. Dong, *An improved upper bound for colorings without symmetrically colored k -term arithmetic progressions*, arXiv:2607.20752, version 2. <https://arxiv.org/abs/2607.20752v2> (sources 09 and 10; source 10 points to its Sections 3–4 and Remark 4.2). Source 09 describes version 2 as “version dated 24 August 2026”. *[write, 6 October 2026; sources 06–10] Corrected:* the arXiv listing, read at the second write, gives version 1 on 22 July 2026 and version 2 on 28 July 2026 (as the batch-116 placement 6f530c29c also records).
- [34] R. Raghavan, *Improved bounds for 3-progressions*, preprint, 2026, version 3, 15 May 2026. [arXiv:2603.27045](https://arxiv.org/abs/2603.27045) (source 10).
- [35] W. T. Gowers, *A uniform set with fewer than expected arithmetic progressions of length 4*, preprint made publicly available in 2020. [arXiv:2004.07598](https://arxiv.org/abs/2004.07598) (source 10).
- [36] G. Petridis, *New proofs of Plünnecke-type estimates for product sets in groups*, Combinatorica **32** (2012), 721–733. [arXiv:1101.3507](https://arxiv.org/abs/1101.3507) (source 10).
- [37] M. Deng, J. Tidor and Y. Zhao, *Uniform sets with few progressions via colorings*, Mathematical Proceedings of the Cambridge Philosophical Society **179** (2025), 79–103. [doi:10.1017/S0305004125000106](https://doi.org/10.1017/S0305004125000106); [arXiv:2307.06914](https://arxiv.org/abs/2307.06914) (source 10).
- [38] J. Fühner and V. Taranchuk, *Large line-free sets and their applications*, preprint, 2024, revised version 2, 2026. [arXiv:2403.18611](https://arxiv.org/abs/2403.18611), Section 6 (source 10).
- [39] C. Chevalley, *Démonstration d’une hypothèse de M. Artin*, Abhandlungen aus dem Mathematischen Seminar der Universität Hamburg **11** (1935), 73–75. [doi:10.1007/BF02940714](https://doi.org/10.1007/BF02940714) (source 10).
- [40] E. Warning, *Bemerkung zur vorstehenden Arbeit*, Abhandlungen aus dem Mathematischen Seminar der Universität Hamburg **11** (1935), 76–83. [doi:10.1007/BF02940715](https://doi.org/10.1007/BF02940715) (source 10).
- [41] [source 06] V. Reshetnikov, *ProveIt: Gowers–Szemerédi paper and Lean development*, repository snapshot 7d99e5deeb209a4671060990cb1ea3d672d57064, inspected 6 October 2026. [Edited paper](#); [Lean sources](#); [formalization checkpoint](#); [statement ledger](#).
- [42] [source 07] V. Reshetnikov, ProveIt repository, *GowersSzemerédi/Section10.lean*, at commit ae0e895748209f0dd079c6b5ca6df6c6fc065a81, inspected 6 October 2026. Source blob 302e8a223f56dcdabf29f30ca3c80ac62a7adbf.c. [Pinned statement source](#). The repository’s paper transcription is under `Combinatorics/Ramsey/Papers/sz-thm-gowers-proof/`.

- [43] [source 07] V. Reshetnikov, ProveIt repository, *GowersSzemerédi/Proofs10Shift.lean*, at the same pinned commit. Source blob `ab399d43dce38f7c8387cd011e2df3915a75ee19`. [Pinned proof source](#), including `lemma_10_9_holds` and its nonvacuous two-step preparation.
- [44] [source 07] V. Reshetnikov, ProveIt repository, *GowersSzemerédi/Audit.lean*, inspected 6 October 2026. Source blob `10b6a30720f41111569bc147d693d7f86508dab0`. [Audit source](#). Its distinction between an axiom audit and completion of conditional mathematical inputs is relevant to the status of source 07.
- [45] [source 08] V. Reshetnikov, ProveIt repository, *Definitions.lean*, GowersSzemerédi development, *Combinatorics/Ramsey/Lean*. Read-only source inspection, 6 October 2026. [Repository source](#) (a link to the moving branch `main`; the blob identifier in source 08’s `source_provenance.json` is the permanent record).
- [46] [source 08] V. Reshetnikov, ProveIt repository, *Proofs03CubeUpper.lean*, GowersSzemerédi development, *Combinatorics/Ramsey/Lean*. Read-only source inspection, 6 October 2026. [Repository source](#) (moving branch `main`; see the previous entry).
- [47] [source 09] V. Reshetnikov, *ProveIt*, GitHub repository, *Combinatorics/Ramsey* source paper and Lean development. Inspected checkpoint `7d99e5deeb209a4671060990cb1ea3d672d57064`, 6 October 2026. <https://github.com/VladimirReshetnikov/ProveIt>. Source 09’s source audit (shipped as `09-cubical-expansions-SOURCE_AUDIT.md`) records the relevant paths and blob identifiers.
- [48] [source 10] V. Reshetnikov, *ProveIt*, repository snapshot `7d99e5deeb209a4671060990cb1ea3d672d57064`, 6 October 2026. [Edited paper](#); [Lean development](#); [source status ledger](#).

Added in the third write (sources 11–14, 6 October 2026). Works above that sources 11–14 also cite: [1] (all four; source 11: “The original page 500 was checked against the published PDF”, citing pp. 498–501; source 12: “Main comparison: Lemma 10.10, Corollary 10.11, Lemma 10.12, Theorem 10.13, and Corollary 10.14”; source 13 also gives the author paper at <https://sites.math.rutgers.edu/~zeilberg/akherim/GowersMasterpiece.pdf>; source 14: “The supplied repository contains the paper source”); [2] (all four; source 13 as “version 2 (2024)”); [4] (12, as the published paper, [doi:10.1017/S0013091505000325](https://doi.org/10.1017/S0013091505000325): “See Definition 2.6, Lemma 8.2, and the comparison with Gowers’s Lemma 10.10 on p. 84”); [7] (13, arXiv version 2 (2011); 14, as *Journal d’Analyse Mathématique* **117** (2012), 133–186); [29] (13); [40] (14, title “Bemerkung zur vorstehenden Arbeit von Herrn Chevalley”); [33] (14: “version 2, revised 28 July 2026”). The entries below are new.

- [49] M. Blum, M. Luby and R. Rubinfeld, *Self-testing/correcting with applications to numerical problems*, *Journal of Computer and System Sciences* **47** (1993), 549–595. [doi:10.1016/0022-0000\(93\)90044-W](https://doi.org/10.1016/0022-0000(93)90044-W) (source 11).
- [50] S. Fu, X. Feng, Q. Wang and C. Carlet, *On the derivative imbalance and ambiguity of functions*, arXiv:1710.07765, version 2, 7 September 2018. <https://arxiv.org/abs/1710.07765> (source 11).
- [51] D. Conlon and W. T. Gowers, *Freiman homomorphisms on sparse random sets*, *Quarterly Journal of Mathematics* **68** (2017), 275–300. [doi:10.1093/qmath/haw058](https://doi.org/10.1093/qmath/haw058); [arXiv:1603.01734](https://arxiv.org/abs/1603.01734). See the introduction and Section 6 for induced maps and concentrated distributions (source 12).
- [52] Y. Zhao, *Graph Theory and Additive Combinatorics*, MIT OpenCourseWare 18.225, Fall 2023, lectures 19–21: Forbidding 3-Term Arithmetic Progressions; especially the Fourier and Dirichlet approximation sections. https://ocw.mit.edu/courses/18-225-graph-theory-and-additive-combinatorics-fall-2023/mit18_225_f23_lec19-21.pdf (source 13).
- [53] I. P. F. da Silva, *Las Vergnas Cube Conjecture and reconstruction properties of the cube matroid*, arXiv:math/0508103 (2005). <https://arxiv.org/abs/math/0508103> (source 13).
- [54] T. Sanders, *Approximate (Abelian) groups*, in *European Congress of Mathematics*, European Mathematical Society, Zürich, 2013, 675–689. <https://arxiv.org/abs/1212.0456> (source 14).

- [55] T. Karam, *On entropy Marton-type inequalities and small symmetric differences with cosets of abelian groups*, 2024 preprint. <https://arxiv.org/abs/2406.10872> (source 14).
- [56] S. Asgarli, *A new proof of Warning’s second theorem*, The American Mathematical Monthly **125** (2018), 549–553. doi:10.1080/00029890.2018.1449555; <https://webpages.scu.edu/ftp/sasgarli/WST.pdf> (source 14).
- [57] É. Leducq, *A new proof of Delsarte, Goethals and Mac Williams theorem on minimal weight codewords of generalized Reed–Muller code*, Finite Fields and Their Applications **18** (2012), 581–586. <https://arxiv.org/abs/1001.2554> (source 14).
- [58] [source 14] *Sharper Local Bounds in Gowers’s Proof of Szemerédi’s Theorem*, earlier unpublished project report prepared in this collaboration, 6 October 2026. Used to identify the earlier cube coefficient gap; all necessary mathematical inputs are reproved here. [write] This is source 04 of this report.
- [59] [source 14] *Quantitative Refinements Around Gowers’ Proof of Szemerédi’s Theorem*, earlier unpublished project report prepared in this collaboration, 6 October 2026. Contains the heterogeneous-degree polynomial-coloring dimension threshold; the present report reproves it and determines the exact supersaturation count. [write] This is source 10 of this report.
- [60] [source 11] V. Reshetnikov, ProveIt repository, *Gowers paper transcription*, Combinatorics/Ramsey/Papers/sz-thm-gowers-proof/sz-thm-gowers-proof.tex. Commit 40f885098a74607b423909dbcb8f070993b7c2a8; accessed 6 October 2026. [Pinned source](#).
- [61] [source 11] V. Reshetnikov, ProveIt repository, *Sections06_07.lean*, Combinatorics/Ramsey/Lean/GowersSzemerédi/Sections06_07.lean. Same commit and access date. [Pinned statement source](#).
- [62] [source 11] V. Reshetnikov, ProveIt repository, *Proofs05_10.lean*, Combinatorics/Ramsey/Lean/GowersSzemerédi/Proofs05_10.lean. Same commit and access date. [Pinned proof source](#).
- [63] [source 12] V. Reshetnikov, *ProveIt: reconstructed and corrected Gowers source*, repository snapshot of 6 October 2026, commit 9ea383e9b28d885a936d5d681cd5584fffd2cb29. [Combinatorics/Ramsey/Papers/sz-thm-gowers-proof/](#).
- [64] [source 12] V. Reshetnikov, *ProveIt: GowersSzemerédi Section 10 statements and proofs*, same snapshot. [Section10.lean](#); [Proofs10Main.lean](#); [Proofs10Bohr.lean](#); [Proofs10Induced.lean](#).
- [65] [source 12] V. Reshetnikov, *ProveIt: corrected extraction and shifting arguments*, same snapshot. [Proofs10Extraction.lean](#); [Proofs10Shift.lean](#); [Proofs10Regular.lean](#).
- [66] [source 12] V. Reshetnikov, *ProveIt: Gowers Szemerédi formalization checkpoint*, 6 October 2026, same snapshot. [FORMALIZATION_STATUS.txt](#). Build and audit figures are reported by this checkpoint, not independently rerun for this article.
- [67] [source 13] V. Reshetnikov, *ProveIt: adapted Gowers paper source*, commit 9ea383e9b28d885a936d5d681cd5584fffd2cb29, inspected 6 October 2026. <https://github.com/VladimirReshetnikov/ProveIt/blob/9ea383e9b28d885a936d5d681cd5584fffd2cb29/Combinatorics/Ramsey/Papers/sz-thm-gowers-proof/sz-thm-gowers-proof.tex>.
- [68] [source 13] V. Reshetnikov, *ProveIt: GowersSzemerédi Lean development*, same commit. <https://github.com/VladimirReshetnikov/ProveIt/tree/9ea383e9b28d885a936d5d681cd5584fffd2cb29/Combinatorics/Ramsey/Lean/GowersSzemerédi>.
- [69] [source 13] V. Reshetnikov, *ProveIt: Gowers proof source-inventory ledger*, same commit. <https://github.com/VladimirReshetnikov/ProveIt/blob/9ea383e9b28d885a936d5d681cd5584fffd2cb29/Combinatorics/Ramsey/gowers-proof-status.json>.

- [70] [source 14] Vladimir Reshetnikov, *ProveIt: Gowers paper and GowersSzemerédi Lean development*, repository revision 9ea383e9b28d885a936d5d681cd5584fffd2cb29, inspected 6 October 2026. <https://github.com/VladimirReshetnikov/ProveIt/tree/9ea383e9b28d885a936d5d681cd5584fffd2cb29/Combinatorics/Ramsey>. See the accompanying manifest (shipped as `data/14-sharp-extremals-source_manifest.json`) for exact file paths and blob identifiers.
- Added in the fourth write (sources 15–20, 6 October 2026).* Works above that sources 15–20 also cite: [1] (all six; source 15 consulted the public PDF under `users/gasarch` and gives printed pages 559–560 for Lemmas 14.7–14.8; source 16 cites it as “Author paper copy” under `~gasarch`; source 17 as “Primary text consulted through the University of Maryland-hosted PDF and the user-specified repository transcription”, pages 515–518 for Section 9; source 18 as “Published-paper mirror” under `~gasarch`; source 19 gives the PDF at <https://sites.math.rutgers.edu/~zeilberg/akherim/GowersMasterpiece.pdf>; source 20 “Primary paper, especially Section 3”, under `users/gasarch`); [2] (sources 15, 17, 18, 19 and 20; 15, 18, 19 and 20 as version 2); [3] (source 17, as arXiv:2402.17994, 2024); [7] (source 15, as arXiv:1012.3509 version 2, 2011). The entries below are new.
- [71] V. Kovač and K. M. Rogers, *Estimates for L^p variants of Gowers norms*, arXiv:2609.08916v1, 8 September 2026. <https://arxiv.org/abs/2609.08916> (source 15; not checked at the write).
- [72] J. Maynard, *Simultaneous small fractional parts of polynomials*, Geometric and Functional Analysis **31** (2021), 150–179. doi:10.1007/s00039-021-00559-3; <https://arxiv.org/abs/2011.12275> (sources 16, 18 and 19; source 19 cites its Theorem 1.1 and capitalizes the title). The journal data are as the sources give them; the intake confirmed the arXiv record, not the journal reference.
- [73] C. F. (Joshua) Lau, *Simultaneously small fractional parts of polynomials*, arXiv:2407.01611v1, 27 June 2024, preprint; Theorem 1.1 and Corollary 1.2. <https://arxiv.org/abs/2407.01611> (sources 16, 18 and 19; source 18 uses Corollary 1.2, source 19 Theorem 1.1 as an optional numerical input). *[write]* The intake’s check of arXiv lists the author as Cheuk Fung Lau; the sources’ “(Joshua)” is kept as delivered.
- [74] N. Beerenwinkel, L. Pachter and B. Sturmfels, *Epistasis and shapes of fitness landscapes*, Statistica Sinica **17** (2007), 1317–1342. <https://www3.stat.sinica.edu.tw/statistica/oldpdf/A17n43.pdf>. See especially the circuit classification in Example 3.9 (source 18).
- [75] D. Berend and A. Kontorovich, *The missing mass problem*, arXiv:1111.2328 (2011). <https://arxiv.org/abs/1111.2328> (source 19).
- [76] G. Decrouez, M. Grabchak and Q. Paris, *Finite sample properties of the mean occupancy counts and probabilities*, arXiv:1601.06537v2 (2016). <https://arxiv.org/abs/1601.06537> (source 19).
- [77] T. Tao and T. Ziegler, *Concatenation theorems for anti-Gowers-uniform functions and Host–Kra characteristic factors*, Discrete Analysis 2016:13, 60 pp. Version 4, 17 January 2017. <https://arxiv.org/abs/1603.07815>; DOI 10.19086/da.850 (source 20).
- [78] [source 15] V. Reshetnikov and repository contributors, *ProveIt: Gowers Szemerédi development*, commit 9ea383e9b28d, accessed 6 October 2026. [Pinned repository source](#).
- [79] [source 15] ProveIt, `Sections14_15.lean`, especially the product property, arrangement definitions, and repaired Corollary 14.6 and Lemma 14.8, at the commit in [78]. [Pinned repository source](#).
- [80] [source 15] ProveIt, `Proofs14HigherArrangements.lean`, general even-moment identity, second-moment estimate, interpolation, and `lemma_14_7_holds`, at the commit in [78]. [Pinned repository source](#).
- [81] [source 15] ProveIt, `Proofs14Arrangements.lean`, reference-cube proof of Corollary 14.6, at the commit in [78]. [Pinned repository source](#).

- [82] [source 15] ProveIt, `Proofs14ProductToArrangements.lean`, proof of the corrected Lemma 14.8, at the commit in [78]. [Pinned repository source](#).
- [83] [source 15] ProveIt, `FORMALIZATION_STATUS.txt`, checkpoint dated 6 October 2026, at the commit in [78]. [Pinned repository source](#).
- [84] [source 16] Vladimir Reshetnikov, *ProveIt: Gowers Szemerédi paper transcription, Lean development, formalization status, and local quantitative refinements*. Snapshot `dca9f06383ccdfc2ed208ccb7c988221fdb09ac`, inspected 6 October 2026. <https://github.com/VladimirReshetnikov/ProveIt/tree/dca9f06383ccdfc2ed208ccb7c988221fdb09ac/Combinatorics/Ramsey>. [write] The “local quantitative refinements” at this snapshot are this report’s directory with source 04’s manuscript staged as `article.tex` (Section F.4).
- [85] [source 17] V. Reshetnikov, *ProveIt*, `Combinatorics/Ramsey/Lean/GowersSzemerédi/Sections08_09.lean`, commit `dca9f06383ccdfc2ed208ccb7c988221fdb09ac`. [Pinned source](#). Accessed 6 October 2026.
- [86] [source 17] V. Reshetnikov, *ProveIt*, `Proofs09Moments.lean`, same directory and commit as above. [Pinned source](#). In particular, the inspected range includes `section9_second_moment_1e`, the exact equality $N^2|B|$ inside its proof, and `lemma_9_2_holds`.
- [87] [source 17] V. Reshetnikov, *ProveIt*, `Proofs09Restriction.lean`, same directory and snapshot. [Pinned source](#). The inspected source includes the height-zero arrangement equivalences and imports of the moment and general restriction modules.
- [88] [source 18] V. Reshetnikov, *ProveIt: Gowers–Szemerédi paper reconstruction and Lean development*, repository snapshot inspected 6 October 2026, commit `dca9f06383ccdfc2ed208ccb7c988221fdb09ac`. <https://github.com/VladimirReshetnikov/ProveIt/tree/dca9f06383ccdfc2ed208ccb7c988221fdb09ac/Combinatorics/Ramsey>.
- [89] [source 19] V. Reshetnikov and contributors, *ProveIt: edited transcription of Gowers’ paper*, snapshot `dca9f06383ccdfc2ed208ccb7c988221fdb09ac`, inspected 6 October 2026. <https://github.com/VladimirReshetnikov/ProveIt/blob/dca9f06383ccdfc2ed208ccb7c988221fdb09ac/Combinatorics/Ramsey/Papers/sz-thm-gowers-proof/sz-thm-gowers-proof.tex>.
- [90] [source 19] V. Reshetnikov and contributors, *ProveIt: GowersSzemerédi Lean development*, same snapshot. <https://github.com/VladimirReshetnikov/ProveIt/tree/dca9f06383ccdfc2ed208ccb7c988221fdb09ac/Combinatorics/Ramsey/Lean/GowersSzemerédi>. The accompanying source manifest records the relevant individual files and their hashes.
- [91] [source 19] V. Reshetnikov and contributors, *ProveIt: numbered-statement ledger and formalization checkpoint*, same snapshot. [Statement ledger](#); [checkpoint report](#).
- [92] [source 19] V. Reshetnikov and contributors, *ProveIt: GowersSzemerédi axiom audit*, same snapshot. <https://github.com/VladimirReshetnikov/ProveIt/blob/dca9f06383ccdfc2ed208ccb7c988221fdb09ac/Combinatorics/Ramsey/Lean/GowersSzemerédi/Audit.lean>.
- [93] [source 20] ProveIt contributors, *Sharper Local Bounds in Gowers’s Proof of Szemerédi’s Theorem*, inspected research manuscript in `Combinatorics/Ramsey/Research/GowersSzemerédi/local-quantitative-refinements/article.tex`. Inspected blob `45b408f6d970a6d997973189663395670b5f3e6e`; 6 October 2026. Repository: <https://github.com/VladimirReshetnikov/ProveIt/tree/main/Combinatorics/Ramsey>. Authorship of individual contributed research reports is not inferred from repository ownership. [write] This blob is source 04’s delivered manuscript as staged by `aefb0b44a`; source 20 never saw the merged report, and in particular not source 06’s quadratic counting theorem (note in Section F.25).

- [94] [source 20] ProveIt contributors, `GowersSzemerédi/Sections01_03.lean`, under `Combinatorics/Ramsey/Lean`, at commit `1bd57e91c76800c175c311fd1f37b2a10f07b9db`. Source specifications for the Section 3 statements; inspected 6 October 2026.
- [95] [source 20] ProveIt contributors, `GowersSzemerédi/Proofs01_03.lean`, inspected blob `e72f7bd53a23d5880cbddbfebd0664d529e47c10`. Finite Fourier identities in the namespace `LeanProofs.GowersSzemerédi`; inspected 6 October 2026.
- Added in the fifth write (sources 21–27, 6 October 2026).* Works above that sources 21–27 also cite: [1] (all seven; source 21 gives the original-paper copy under `~gasarch` and “Lemmas 3.8–3.10, printed pp. 483–485”; source 22 gives the DOI and the copy under `~gasarch`; source 23 likewise; source 24 “Section 7, especially Proposition 7.3, Lemmas 7.4–7.5, and Corollary 7.6”, inspected under `users/gasarch`; source 25 “Sections 14–15, especially pp. 562–565”, copy under `users/gasarch`; source 26 as “Author paper mirror”, “Section 5, particularly Lemmas 5.4, 5.12 and 5.13”; source 27 “Section 15, especially Lemmas 15.1–15.5, pp. 561–565”); [2] (sources 21, 23 and 24, each as version 2; source 21 cites its Theorem 1.1, source 24 “only to place the local improvements in a modern global context”); [7] (source 21, as arXiv:1012.3509 (2010), “subsequent published version in Analysis & PDE (2012)”); [31] (source 21); [11] (source 24, as version 2 (2010), “Background reference for the established sampling method, not an external dependency of the proofs given here”); [36] (source 24, with the issue number 6 and arXiv version 3); [12] (source 24, with arXiv version 2: “Theorem 1.2 is the explicitly marked external input in Corollary II.80.2; Lemma 2.1 is a related one-neighbourhood selection argument”); [40] (source 22, under the full title *Bemerkung zur vorstehenden Arbeit von Herrn Chevalley*); [56] and [57] (source 22); [33] (source 22, as a “2026 preprint, arXiv:2607.20752, version 2”: “The field-norm and layered constructions are relevant prior work; their method is not claimed as new here”). The entries below are new.
- [96] P. Delsarte, J.-M. Goethals and F. J. MacWilliams, *On generalized Reed–Muller codes and their relatives*, Information and Control **16** (1970), 403–442. doi:10.1016/S0019-9958(70)90214-7 (source 22).
- [97] J. B. Roberts, *Note on linear forms*, Proceedings of the American Mathematical Society **7** (1956), 465–469. doi:10.1090/S0002-9939-1956-0091961-5. The conductor formula used here is proved directly in Proposition I.37.3 (source 23).
- [98] C. A. Athanasiadis, *Characteristic polynomials of subspace arrangements and finite fields*, Advances in Mathematics **122** (1996), 193–233. <https://www.sciencedirect.com/science/article/pii/S0001870896900596>. This is background for the further question on intersection lattices; no theorem from it is needed for the elementary bounds proved above (source 25).
- [99] R. P. Stanley, *An introduction to hyperplane arrangements*, in *Geometric Combinatorics*, E. Miller, V. Reiner and B. Sturmfels, eds., IAS/Park City Mathematics Series **13**, American Mathematical Society, 2007, pp. 389–496. Author’s notes and errata: <https://math.mit.edu/~rstan/arrangements/arr.html>. Background only; the hyperplane counting facts used here are proved in the text (source 27).
- [100] Y. Zhao, *Graph Theory and Additive Combinatorics: Exploring Structure and Randomness*, Cambridge University Press, 2023. Author version, updated 18 June 2024, Lemmas 6.4.5–6.4.6, printed pages 219–220 (source 26). [write] Not the lecture notes [52] of source 13.
- [101] M. Lewko, *A Fourier-Free Density-Increment Proof of Roth’s Theorem*, arXiv:2605.19310v1, 19 May 2026; Sections 5–6. Cited for method context, not as a dependency of the present proofs (source 26).
- [102] [source 21] ChatGPT, prepared for Vladimir Reshetnikov, *Sharp Extremal Refinements in Gowers’s Szemerédi Framework: Secondary Cube Asymptotics, Arrangement Stability, and Polynomial Supersaturation*, unpublished project report, 6 October 2026. Source file: `gowers_extremal_refinements.tex`. Section 3 proves the exponent-five correction exponent and

- leaves its optimal coefficient undetermined. This source is credited as prior project work, not as a refereed publication. *[write]* This is source 14 of this report; its Section 3 is Section III.35.
- [103] [source 21] VladimirReshetnikov/ProveIt, *Combinatorics/Ramsey*, repository snapshot 43306653afaaaaa0b58a7ca71e4067c245594867, inspected 6 October 2026. <https://github.com/VladimirReshetnikov/ProveIt/tree/43306653afaaaaa0b58a7ca71e4067c245594867/Combinatorics/Ramsey>. Research compilation: *Research/GowersSzemerédi/local-quantitative-refinements/article.tex*. *[write]* At this snapshot the compilation was the first write (sources 01–05).
- [104] [source 22] Vladimir Reshetnikov, *ProveIt: Ramsey papers, research, and GowersSzemerédi formalization*, repository snapshot 43306653afaaaaa0b58a7ca71e4067c245594867, inspected 6 October 2026. <https://github.com/VladimirReshetnikov/ProveIt/tree/43306653afaaaaa0b58a7ca71e4067c245594867/Combinatorics/Ramsey>. The transcription contains editorial repairs and should be distinguished from the printed paper when quoting a quantitative statement.
- [105] [source 22] *Higher-Order Cube Expansions and Sharp Degeneracy Bounds in Gowers’ Szemerédi Framework: Simultaneous polynomial partitions, exact constants, and quantitative obstructions*, earlier unpublished project report prepared for ProveIt, 6 October 2026; source file *gowers_cubes_degeneracy_partitions.tex*. Used for the previously established $d \geq 2$ first-order coefficient and the paired-moment rank argument, which are reproved here. *[write]* This is source 18 of this report.
- [106] [source 22] *Sharp Extremal Refinements in Gowers’s Szemerédi Framework: Secondary cube asymptotics, arrangement stability, and polynomial supersaturation*, earlier unpublished project report prepared for ProveIt, 6 October 2026; source file *gowers_extremal_refinements.tex*. Used for the previous low-degree symmetric supersaturation result and the explicit global quintic question. Its necessary ingredients are reproved here. *[write]* This is source 14 of this report.
- [107] [source 23] V. Reshetnikov, *ProveIt repository: revised transcription of Gowers’s paper*, revision 65f4e039fc4ad30a8f492e63abde2bcd0740d7c9, consulted 6 October 2026. <https://github.com/VladimirReshetnikov/ProveIt/blob/65f4e039fc4ad30a8f492e63abde2bcd0740d7c9/Combinatorics/Ramsey/Papers/sz-thm-gowers-proof/sz-thm-gowers-proof.tex>.
- [108] [source 23] ProveIt research collection, *Local quantitative refinements of Gowers’s proof of Szemerédi’s theorem*, revision 65f4e039fc4ad30a8f492e63abde2bcd0740d7c9, consulted 6 October 2026. <https://github.com/VladimirReshetnikov/ProveIt/blob/65f4e039fc4ad30a8f492e63abde2bcd0740d7c9/Combinatorics/Ramsey/Research/GowersSzemerédi/local-quantitative-refinements/article.tex>. *[write]* At this revision the article was the first write (sources 01–05).
- [109] [source 23] ProveIt, *Gowers Szemerédi formalization checkpoint*, 6 October 2026, revision 65f4e039fc4ad30a8f492e63abde2bcd0740d7c9. https://github.com/VladimirReshetnikov/ProveIt/blob/65f4e039fc4ad30a8f492e63abde2bcd0740d7c9/Combinatorics/Ramsey/FORMALIZATION_STATUS.txt.
- [110] [source 23] ProveIt, *GowersSzemerédi Lean development*, revision 65f4e039fc4ad30a8f492e63abde2bcd0740d7c9. Relevant modules: *Proofs16CommonBase.lean*, *Proofs16RelativeSteps.lean*, *Proofs16Interpolation.lean*, and *Proofs16SampleSelection.lean*. <https://github.com/VladimirReshetnikov/ProveIt/tree/65f4e039fc4ad30a8f492e63abde2bcd0740d7c9/Combinatorics/Ramsey/Lean/GowersSzemerédi>.
- [111] [source 23] ProveIt project report, *Sharp Local Refinements in Gowers’ Szemerédi Framework*, 6 October 2026, unpublished research draft. Theorem 3.8, Corollary 3.9, and Research Question 11.1. Previously delivered as *gowers_local_refinements(1).pdf*; the report’s source revision is 9ea383e9b28d885a936d5d681cd5584fffd2cb29. *[write]* This is source 13 of this report: its Theorem 3.8 and Corollary 3.9 are Theorem I.26.8 and Corollary I.26.9, and its Research Question 11.1 is Research question V.54.1 (“13 Question 1”, item I.12).

- [112] [source 23] ProveIt project report, *Sharp Extremal Refinements in Gowers’s Szemerédi Framework: Secondary Cube Asymptotics, Arrangement Stability, and Polynomial Supersaturation*, 6 October 2026, unpublished research draft. Previously delivered as `gowers_extremal_refinements.pdf`. The order-five secondary coefficient was exhibited there as a lower example, with its optimality left open. *[write]* This is source 14 of this report.
- [113] [source 24] V. Reshetnikov, *ProveIt*, GitHub repository. Snapshot inspected: 6 October 2026, commit `2e0553b0c8a6c9c93221156cc83e95d68ee8bb07`. <https://github.com/VladimirReshetnikov/ProveIt>. Relevant paths under `Combinatorics/Ramsey/`: `Papers/sz-thm-gowers-proof/sz-thm-gowers-proof.tex`; `Lean/GowersSzemerédi/Sections06_07.lean`; `Lean/GowersSzemerédi/Proofs07DRC.lean`; `Lean/GowersSzemerédi/Proofs07BalogSzemerédi.lean`. The accompanying provenance note records immutable links (shipped as `24-section7-compression-PROVENANCE.md`).
- [114] [source 25] Vladimir Reshetnikov, *ProveIt*, `Combinatorics/Ramsey/Lean/GowersSzemerédi/Sections14_15.lean`, commit `568d1802e99404c0710be33cceb401d25211f445`, blob `5d91dce37bbca59fc76dfd30ec3a485d72ba583c`. Consulted 6 October 2026. https://github.com/VladimirReshetnikov/ProveIt/blob/568d1802e99404c0710be33cceb401d25211f445/Combinatorics/Ramsey/Lean/GowersSzemerédi/Sections14_15.lean.
- [115] [source 25] Vladimir Reshetnikov, *ProveIt: Local Quantitative Refinements of Gowers’s Proof of Szemerédi’s Theorem*, research README, observed on `main`, 6 October 2026. This reference is a scope audit, not a mathematical premise. <https://github.com/VladimirReshetnikov/ProveIt/blob/main/Combinatorics/Ramsey/Research/GowersSzemerédi/local-quantitative-refinements/README.md>. *[write]* When source 25 was written (pin `568d1802e`, 14:52) the README was that of the first write (sources 01–05); source 18’s manuscript was not in the repository.
- [116] [source 26] V. Reshetnikov, *ProveIt*, pinned [Section05.lean](#), [Corollary 5.8](#), commit `c94998a5db18ba63c815a3bf6c3877a0fde91930`.
- [117] [source 26] V. Reshetnikov, *ProveIt*, [Definitions.lean](#), same pin, polynomial and progression definitions.
- [118] [source 26] V. Reshetnikov, *ProveIt*, [Proofs05Corollary58.lean](#), same pin, common-target scale repair.
- [119] [source 26] V. Reshetnikov, *ProveIt*, [FORMALIZATION_STATUS.txt](#), same pin, Section 16 source-fidelity qualifications. Repository status is cited as recorded, not as a build rerun here.
- [120] [source 27] V. Reshetnikov, repository maintainer, *ProveIt: Combinatorics/Ramsey*, repository materials inspected 6 October 2026. <https://github.com/VladimirReshetnikov/ProveIt/tree/main/Combinatorics/Ramsey>. In particular, `Lean/GowersSzemerédi/Sections14_15.lean`, blob `5d91dce37bbca59fc76dfd30ec3a485d72ba583c`; and the edited Gowers transcription, blob `41085a0efde8c1932e86e80c791984f52841c22e`. These identifiers describe inspected source snapshots, not formal verification of this article. *[write]* Both blobs are unchanged at HEAD.

Added in the sixth write (sources 28–33, 6 October 2026). Works above that sources 28–33 also cite: [1] (all six; source 28 with the DOI, “The Section 10 statements are also available in the repository’s transcription”; source 29 the copy under `~gasarch`; source 30 “Lemma 10.10 and Corollary 10.11, printed pp. 529–530, and Lemma 10.12, pp. 530–531”; source 31 the DOI and the copy under `~gasarch`; source 32 “Section 5, especially printed pp. 490–493” through a copy under `~zeilberg`; source 33 the DOI and the copy under `users/gasarch`); [2] (sources 29 and 33, as version 2 or without version); [3] (source 28, as version 3, 10 August 2024); [7] (sources 28, 29 and 31; source 31 as Journal d’Analyse Mathématique **117** (2012), 133–186, “Theorem 1.4, with $k = 3$ and $p_k = 2$ ”; source 28 as version 2, 26 June 2011); [49] (sources 29 and 33; source 33 adds the issue number 3 and an author-hosted copy);

- [4] (source 31, “proof of Lemma 6.2”); [12] (source 31, with arXiv version 2 of 18 February 2024: “The explicit input used here is Theorem 1.2”); [36] (source 31); [72] (source 32, as arXiv:2011.12275v1, 24 November 2020, “Theorem 1.1 and Corollary 1.2”). The entries below are new.
- [121] S. Foucart and H. Rauhut, *A Mathematical Introduction to Compressive Sensing*, Applied and Numerical Harmonic Analysis, Birkhäuser/Springer, 2013. Section 8.3, proof of Theorem 8.9; accessible manuscript: https://users.math.msu.edu/users/iwenmark/Teaching/MTH994/Holger_Simon_book.pdf (source 29).
- [122] M. Ben-Or, D. Coppersmith, M. Luby and R. Rubinfeld, *Non-Abelian homomorphism testing, and distributions close to their self-convolutions*, author manuscript. <https://people.csail.mit.edu/ronitt/papers/convd.pdf> (source 29).
- [123] F. Nicola, *Maximally localized Gabor orthonormal bases on locally compact Abelian groups*, Advances in Mathematics **451** (2024), article 109786. <https://arxiv.org/abs/2305.02738>; <https://doi.org/10.1016/j.aim.2024.109786> (source 29).
- [124] B. Green, *A Szemerédi-type regularity lemma in abelian groups, with applications*, preprint arXiv:math/0310476, version 2, 21 October 2004. <https://arxiv.org/abs/math/0310476> (source 30, “Sections 3–4”).
- [125] D. Gross, S. Nezami and M. Walter, *Schur–Weyl duality for the Clifford group with applications: property testing, a robust Hudson theorem, and de Finetti representations*, Communications in Mathematical Physics **385** (2021), 1325–1393. <https://arxiv.org/abs/1712.08628> (source 31).
- [126] T. Tao, *254B, Notes 4: The Bourgain–Gamburd expansion machine*, 13 January 2012, Lemma 1 and Exercise 1. <https://terrytao.wordpress.com/2012/01/13/254b-notes-4-the-bourgain-gamburd-expansion-machine/> (source 31).
- [127] E. Breuillard, B. Green, R. Guralnick and T. Tao, *Expansion in finite simple groups of Lie type*, Journal of the European Mathematical Society **17** (2015), 1367–1434, Appendix A. <https://www.imo.universite-paris-saclay.fr/~emmanuel.breuillard/expansion.pdf> (source 31).
- [128] B. Green and T. Tao, *New bounds for Szemerédi’s theorem, II: A new bound for $r_4(N)$* , arXiv:math/0610604v2, 4 January 2024. Section 6 and Appendix A, Proposition A.2. [Primary article PDF](#) (source 32).
- [129] W. T. Gowers and L. Milićević, *An inverse theorem for Freiman multi-homomorphisms*, arXiv:2002.11667 (2020), version 3, 6 September 2021. <https://arxiv.org/abs/2002.11667> (source 33).
- [130] L. Milićević, *An inverse theorem for certain directional Gowers uniformity norms*, Publications de l’Institut Mathématique **113**(127) (2023), 1–56. [doi:10.2298/PIM2327001M](https://doi.org/10.2298/PIM2327001M). Publication metadata also listed on the author’s papers page: <https://www.mi.sanu.ac.rs/~luka.milicevic/papers.html> (source 33). *[write]* Not the 2026 preprint [32] of source 09.
- [131] [source 28] V. Reshetnikov, *ProveIt: Combinatorics/Ramsey*, repository snapshot [e3125e6982abe2d44bd7cb9369063d4e69d41a2d](https://github.com/Reshetnikov/ProveIt/commit/e3125e6982abe2d44bd7cb9369063d4e69d41a2d), 6 October 2026. [Pinned repository tree](#). Consult the accompanying source-provenance record for the inspected manuscripts and the formalization status ledger (shipped as `data/28-log-selection-source_provenance.json`).
- [132] [source 28] ProveIt research collection, source 04, *Sharper Local Bounds in Gowers’s Proof of Szemerédi’s Theorem: polynomial homomorphism restriction, sharp cube estimates, and floor-pattern complexity*, research report, 2026. Included in the combined article at the pinned repository snapshot [131]. Source 04’s own reference pin is 327773149. *[write]* Source 04 of this report, its base.

- [133] [source 28] ProveIt research collection, source 07, *Transport-sensitive rigidity in Gowers’s approximate-homomorphism argument*, 6 October 2026. Historical archive `Gowers_Transport_Rigidity.zip`, retrieved at commit 62e21161f. [Historical archive](#). The manuscript’s own reference pin is ae0e895748209f0dd079c6b5ca6df6cfc065a81. *[write]* Source 07 of this report (its chapter is Section II.11). The pin printed here has 39 hexadecimal digits: source 07’s pin is ae0e895748209f0dd079c6b5ca6df6cfc065a81 ([42]); a 6 is missing after 6df6c. Kept as delivered.
- [134] [source 28] ProveIt research collection, source 12, *Sharp Freiman Extension and Polynomial Bohr Radii: explicit refinements of Section 10 of Gowers’s proof of Szemerédi’s theorem*, 6 October 2026. Historical archive `gowers_bohr_extension.zip`, retrieved at commit 08ab4187e. [Historical archive](#). The manuscript’s own reference pin is 9ea383e9b28d885a936d5d681cd5584fffd2cb29. *[write]* Source 12 of this report (its chapter is Section II.33).
- [135] [source 29] ProveIt research collection, *Local Quantitative Refinements of Gowers’s Proof of Szemerédi’s Theorem*, combined project manuscript and accompanying README. Pinned source: <https://github.com/VladimirReshetnikov/ProveIt/tree/e3125e6982abe2d44bd7cb9369063d4e69d41a2d/Combinatorics/Ramsey/Research/GowersSzemerédi/local-quantitative-refinements>. Unrefereed project source; the provenance manifest identifies the inspected files and the additional reports. *[write]* This report, at its first write (sources 01–05).
- [136] [source 29] ProveIt research collection, source 17, *Polynomial-loss restriction in Gowers’s Szemerédi argument* (6 October 2026), especially Section 8 and its second further question. Historical source: https://github.com/VladimirReshetnikov/ProveIt/blob/060528502e7cb7229c7910b35f8c9ad9e61aa023/docs/incoming/Gowers_Polynomial_Restriction_Research.zip. The manuscript within the archive is `gowers_polynomial_restriction/article.tex`. Unrefereed project manuscript. *[write]* Source 17 of this report. The random linear maps are in its Section 7 (Section II.63), not Section 8, in that archive’s manuscript as in this report; its second further question is “Reducing the ambient-size threshold” (Section V.71).
- [137] [source 29] ProveIt research collection, sources 08–09, *Fourth-order stability of additive cube counts: Torsion-sensitive refinements of Gowers’ Lemmas 3.7 and 3.10; and Arithmetic Circuit Expansions for Cubical Counts: A fourth-order strengthening of Gowers’ Lemma 3.10, exact torsion corrections, and a formalization roadmap* (6 October 2026). Historical sources: https://github.com/VladimirReshetnikov/ProveIt/blob/62e21161f47e52f0bb43f5a79d71da7951a48aee/docs/incoming/gowers_cube_refinements.zip; https://github.com/VladimirReshetnikov/ProveIt/blob/62e21161f47e52f0bb43f5a79d71da7951a48aee/docs/incoming/gowers_cubical_expansions.zip. Unrefereed project manuscripts. *[write]* Sources 08 and 09 of this report.
- [138] [source 30] V. Reshetnikov, *ProveIt: Ramsey / GowersSzemerédi*, repository materials inspected 6 October 2026. Pinned statement comparison at commit 8d65d3c9febe4d0e9efe945e11b53af5f19d1435. [Pinned Ramsey source directory](#); detailed provenance in the accompanying source audit (shipped as 30-fixed-radius-bohr-SOURCE_AUDIT.md).
- [139] [source 31] V. Reshetnikov, *ProveIt: Combinatorics/Ramsey*, repository checkpoint e3125e6982abe2d44bd7cb9369063d4e69d41a2d, 6 October 2026. <https://github.com/VladimirReshetnikov/ProveIt/tree/e3125e6982abe2d44bd7cb9369063d4e69d41a2d/Combinatorics/Ramsey>. The corrected transcription and the Lean files named in Section V.122 were inspected at this checkpoint.
- [140] [source 31] *Local Quantitative Refinements of Gowers’s Proof of Szemerédi’s Theorem*, consolidated ProveIt research report, 6 October 2026, at the checkpoint of [139]. `Combinatorics/Ramsey/Research/GowersSzemerédi/local-quantitative-refinements/article.tex`. Relevant parts: weighted derivative energy, exact full-graph rigidity, elementary BSG, and Theme II’s correlation-preservation and near-equality questions. Research-report status, not kernel verification. *[write]* This report at its first write (sources 01–05); the two questions are items II.1 and II.2.

- [141] [source 31] ProveIt incoming reports *Characteristic Five*, *Degree Filtrations in Gowers’ Ramsey Framework*, and *Sharp Thresholds*, archives `gowers_characteristic_five.zip`, `gowers_degree_stratification.zip`, and `gowers_sharp_thresholds.zip`, inspected in `docs/incoming/` at the checkpoint of [139]. These were checked for overlap; none is used as a theorem input here. *[write]* Sources 21, 22 and 23 of this report.
- [142] [source 32] V. Reshetnikov, *ProveIt*, GowersSzemerédi source, commit `c94998a5db18ba63c815a3bf6c3877a0fde91930`. [Corollary 5.8 predicate](#); [polynomial definition](#). The exact source fingerprints are recorded in `SOURCES.md` (shipped as `32-quadratic-increments-SOURCES.md`).
- [143] [source 33] V. Reshetnikov, repository maintainer, *ProveIt: Combinatorics/Ramsey*, corrected Gowers transcription and local-quantitative-refinements research documentation. Consulted 6 October 2026; final crosswalk pin given in Section V.128. <https://github.com/VladimirReshetnikov/ProveIt/tree/9f83dbe0f15205c10b1e6efc6af69b09969f8e03/Combinatorics/Ramsey>. *[write]* At that pin this report was its second write (sources 01–10).
- Added in the seventh write (sources 34–39, 6 October 2026).* Works above that sources 34–39 also cite: [1] (all six; source 34 the copy under `~gasarch`, “The repository’s edited transcription was inspected separately at the pinned commit”; source 35 and source 37 the DOI, source 37 “Section 5, especially Lemmas 5.2–5.5”; source 36 the DOI, “Section 17, especially Proposition 17.2”; source 38 the DOI and a copy under `~zeilberg`; source 39 the copy under `users/gasarch`, “Section 16, especially the proof of Lemma 16.10”); [2] (sources 35, 36 and 38, as arXiv:2402.17995, sources 36 and 38 version 2); [7] (sources 35 and 38; source 38 as version 2, 2011, “Theorem 1.4”); [71] (source 35, as arXiv:2609.08916v1, 8 September 2026; source 35 notes that its Theorem 1.1 assumes a nonnegative function on a group without compact open subgroups); [125] (source 38, with its DOI); [31] (source 38); [33] (source 38, version 2, “the layered norm construction and its consequences for Gowers’ Conjecture 4.2”); [32] (source 36, as arXiv:2609.33733v1, submitted 27 September 2026, “cited for scope comparison, not used as a proof ingredient”); [72] (source 37, as arXiv:2011.12275, 2020); [75] (source 39). The entries below are new.
- [144] N. Alon and V. H. Vū, *Anti-Hadamard matrices, coin weighing, threshold gates, and indecomposable hypergraphs*, Journal of Combinatorial Theory, Series A **79** (1997), 133–160. DOI: [10.1006/jcta.1997.2780](https://doi.org/10.1006/jcta.1997.2780); [author PDF](#), especially Section 2.4 and its final unimodularity remark (source 34).
- [145] R. Craigen, *The range of the determinant function on the set of $n \times n$ $(0, 1)$ -matrices*, Journal of Combinatorial Mathematics and Combinatorial Computing **8** (1990), 161–171. [Publisher PDF](#). The small maximum determinants appear in its table (source 34, which supplies its own prime-specific exclusion certificates).
- [146] R. Shah, *Determinants of binary matrices achieve every integral value up to $\Omega(2^n/n)$* , Linear Algebra and its Applications **645** (2022), 229–236. [arXiv:2006.04701](#) (source 34).
- [147] Y. Shitov, *The range of 0 – 1 determinants is large*, author-uploaded preprint, 14 November 2025. DOI: [10.13140/RG.2.2.32216.35842](https://doi.org/10.13140/RG.2.2.32216.35842); [author’s full-text posting](#). The composition method is credited to this preprint (source 34).
- [148] B. Green and T. Tao, *The quantitative behaviour of polynomial orbits on nilmanifolds*, Annals of Mathematics **175** (2012), 465–540. <https://doi.org/10.4007/annals.2012.175.2.2> (source 35, “used for context only”).
- [149] S. Gurevich and R. Hadani, *Notes on quantization of symplectic vector spaces over finite fields*, [arXiv:0708.0669v3](#) (2009; first version 2007). Background on the finite Heisenberg representation and the Stone–von Neumann theorem (source 36).

- [150] J. Tidor, *Quantitative bounds for the U^4 -inverse theorem over low characteristic finite fields*, Discrete Analysis **2022**:14. doi:10.19086/da.38591; arXiv:2109.13108v2 (source 36).
- [151] L. Milićević, *Approximately symmetric forms far from being exactly symmetric*, arXiv:2112.14755v1 (2021) (source 36).
- [152] Z. Yang and L. Xie, *Frequency-selective Vandermonde decomposition of Toeplitz matrices with applications*, arXiv:1605.02431v3 (2017), accepted by Signal Processing. <https://arxiv.org/abs/1605.02431> (source 37).
- [153] R. Baker, *Small fractional parts of polynomials*, arXiv:1602.04245v2 (2016). <https://arxiv.org/abs/1602.04245> (source 37).
- [154] A. Painsky, *Convergence guarantees for the Good–Turing estimator*, Journal of Machine Learning Research **23**(279) (2022), 1–37. <https://www.jmlr.org/papers/v23/21-1528.html> (source 39).
- [155] [source 34] ProveIt research collection, *Local Quantitative Refinements of Gowers’s Proof of Szemerédi’s Theorem*, merged sources 01–10, 6 October 2026. Pinned TeX source. In particular: source 04, sharp real centered odd-order U^3/U^2 comparison; source 07, transport-sensitive rigidity and its nonuniform-threshold question; source 09, arithmetic circuit expansions and the prime-onset questions. [write] This report at its second write.
- [156] [source 34] Research report prepared for Vladimir Reshetnikov with ChatGPT, *Further Quantitative Refinements of Gowers’s Szemerédi Argument: higher norm comparisons, joint quadratic rigidity, and restriction by affine rank*, 6 October 2026. Pinned incoming archive. The positive autocorrelation method and the bound $c_4^{\text{odd}} \leq (8/35)^{1/4}$ are prior inputs. [write] Source 29 of this report.
- [157] [source 34] Research report prepared for Vladimir Reshetnikov with ChatGPT, *Weighted Derivative Spectra and Sharp Local Rigidity: explicit Freiman restriction, an optimal near-extremal U^3 profile, and the limits of stability for correlation weights*, 6 October 2026. Pinned incoming archive. Inspected as part of the nonduplication comparison. [write] Source 31 of this report; listed in source 34’s bibliography and not cited in its text.
- [158] [source 35] V. Reshetnikov, ProveIt repository, edited transcription of Gowers’ paper, snapshot 12da356730b12cb73c276a258aaadf3223e3662b. File: Combinatorics/Ramsey/Papers/sz-thm-gowers-proof/sz-thm-gowers-proof.tex. Source blob: 41085a0efde8c1932e86e80c791984f52841c22e. <https://github.com/VladimirReshetnikov/ProveIt/tree/12da356730b12cb73c276a258aaadf3223e3662b/Combinatorics/Ramsey>.
- [159] [source 35] ProveIt research collection, *Local Quantitative Refinements of Gowers’s Proof of Szemerédi’s Theorem*, merged research article, 6 October 2026. At the source snapshot in [158]: Combinatorics/Ramsey/Research/GowersSzemerédi/local-quantitative-refinements/article.tex. Source blob: 3adb2722cddc65cb8aecddba96d548021b871cdb. The inspected file prints sources 01–10; later archives were also inspected separately. [write] This report at its second write.
- [160] [source 35] ProveIt incoming research report, *One-stage Fejér selection for Gowers arrangements: single-monomial relation isolation, polynomial parameter loss, and sharp exponents within a positive-Fourier selection class*, 6 October 2026. Archive docs/incoming/fejér-arrangement-selection.zip at the snapshot in [158]; inner manuscript fejér-arrangement-selection/fejér-arrangement-selection.tex. [write] Source 40 of this report (arrival 13f105f42), placed by 2dfe46f44 and not yet written in.
- [161] [source 35] ProveIt research report, source 22, *Degree Filtrations in Gowers’ Ramsey Framework: sharper degeneracy bounds, exact polynomial pattern thresholds, and obstructions to extending norm constructions*, 6 October 2026. Original archive docs/incoming/gowers_degree_stratification.zip, recoverable at arrival commit 26908b3b6; inner manuscript gowers_degree_stratification/gowers_degree_stratification.tex. [write] Source 22 of this report.

- [162] [source 35] ProveIt incoming research report, *Logarithmic Selection, Local Freiman Models, and Higher Uniformity: refinements of Gowers’s Szemerédi framework*, 6 October 2026. Archive docs/incoming/gowers_logarithmic_selection.zip at the snapshot in [158]; inner manuscript gowers_logarithmic_selection/gowers_logarithmic_selection.tex. [write] Source 28 of this report.
- [163] [source 35] ProveIt incoming research report, *Further Quantitative Refinements of Gowers’s Szemerédi Argument: higher norm comparisons, joint quadratic rigidity, and restriction by affine rank*, 6 October 2026. Archive docs/incoming/gowers_moments_stability_restriction.zip at the snapshot in [158]; manuscript gowers_moments_stability_restriction.tex. [write] Source 29 of this report.
- [164] [source 35] ProveIt research report, source 18, *Higher-Order Cube Expansions and Sharp Degeneracy Bounds in Gowers’s Szemerédi Framework: simultaneous polynomial partitions, exact constants, and quantitative obstructions*, 6 October 2026. Original archive docs/incoming/gowers_cubes_degeneracy_partitions.zip, recoverable at commit f7efeb4f5a7250876ac6fed4116d5fb9ebecd543; inner manuscript gowers_cubes_degeneracy_partitions/gowers_cubes_degeneracy_partitions.tex. [write] Source 18 of this report.
- [165] [source 35] ProveIt research report, source 23, *Sharp Thresholds and Extremal Constants in Gowers’s Szemerédi Framework: numerical semigroups, characteristic-five cube asymptotics, and common partitions for affine fibres*, 6 October 2026. Original archive docs/incoming/gowers_sharp_thresholds.zip, recoverable at arrival commit 26908b3b6; inner manuscript gowers_sharp_thresholds/gowers_sharp_thresholds.tex. [write] Source 23 of this report.
- [166] [source 36] V. Reshetnikov, *ProveIt*, Ramsey subtree, snapshot 6c2e2172bab8 (6 October 2026). [Pinned repository tree](#). Corrected paper transcription, source lines 3530–3610; research README and formalization-status file inspected. Full blob identifiers are in SOURCE_AUDIT.md (shipped as 36-symmetry-certification-SOURCE_AUDIT.md).
- [167] [source 37] V. Reshetnikov, maintainer, *ProveIt: Combinatorics/Ramsey*. Inspected Section 5 sources at revision 869289aa38d59354e8dac2213df08ae0dd256415, 6 October 2026. <https://github.com/VladimirReshetnikov/ProveIt/tree/869289aa38d59354e8dac2213df08ae0dd256415/Combinatorics/Ramsey>. See the source audit in Appendix V.141.
- [168] [source 38] V. Reshetnikov and contributors, *ProveIt: Ramsey development*, inspected at commit 12da356730b12cb73c276a258aaadf3223e3662b, 6 October 2026. [Pinned Ramsey directory](#). Exact inspected paths and blob hashes are in SOURCE_LEDGER.json (shipped as data/38-quotients-SOURCE_LEDGER.json).
- [169] [source 38] ProveIt research collection, *Local Quantitative Refinements of Gowers’s Proof of Szemerédi’s Theorem*, consolidated AI-assisted research report, 6 October 2026, at the checkpoint in [168]. The inspected article incorporates source reports 01–10. [Pinned consolidated report](#). [write] This report at its second write.
- [170] [source 38] *Sharp Characteristic-Five Cube Bounds in Gowers’s Framework*, AI-assisted research report prepared for Vladimir Reshetnikov, 6 October 2026. Inspected source: gowers_characteristic_five.tex. This report proves the sharp secondary coefficient uniformly in dimension and the eighth-order coefficient on the five-point group. The source hash is recorded in SOURCE_LEDGER.json. [write] Source 21 of this report.
- [171] [source 38] *Sharp Thresholds and Extremal Constants in Gowers’s Szemerédi Framework*, AI-assisted research report prepared for Vladimir Reshetnikov, 6 October 2026. Inspected source: gowers_sharp_thresholds.tex; see the sharp secondary three-cube asymptotics and further questions. The source hash is recorded in SOURCE_LEDGER.json. [write] Source 23 of this report.

- [172] [source 38] *Weighted Derivative Spectra and Sharp Local Rigidity in Gowers’ Argument*, AI-assisted research report prepared for Vladimir Reshetnikov, 6 October 2026. Inspected source: `gowers_weighted_spectrum_rigidity.tex`. Its coarse subgroup argument and all-group extremal profile precede the sharp odd-order refinement here. The source hash is recorded in `SOURCE_LEDGER.json`. *[write]* Source 31 of this report.
- [173] [source 38] *Degree Filtrations in Gowers’ Ramsey Framework*, AI-assisted research report prepared for Vladimir Reshetnikov, 6 October 2026. Inspected source: `gowers_degree_stratification.tex`. See its exact polynomial-pattern bounds and independent-extension obstruction. The source hash is recorded in `SOURCE_LEDGER.json`. *[write]* Source 22 of this report.
- [174] [source 39] V. Reshetnikov’s `ProveIt` repository, inspected snapshot `869289aa38d59354e8dac2213df08ae0dd256415, Proofs16DistinctSampling.lean`, under `Combinatorics/Ramsey/Lean/GowersSzemerédi/`. Corrected distinct-anchor counts and the repeated-sampling obstruction. [Pinned source at the inspected commit](#).
- [175] [source 39] Same repository and snapshot, `Proofs16AnchoredFibres.lean`. The corrected affine-cover anchored good-set theorem. [Pinned source at the inspected commit](#).
- [176] [source 39] Same repository and snapshot, `Proofs16ShortColumns.lean`. Directly sampled points, short-column recovery, and the `Section16SampledOrAnchoredOn` interface. [Pinned source at the inspected commit](#).
- [177] [source 39] Same repository and snapshot, `Proofs16RecoveredSliceCover.lean`. The existing common slice-cover bounds and candidate indexing. [Pinned source at the inspected commit](#).
- [178] [source 39] Same repository and snapshot, `Proofs16SynchronizedCover.lean`. Geometric synchronization without additional candidates or mass loss. [Pinned source at the inspected commit](#).
- [179] [source 39] Same repository and snapshot, `Proofs16SynchronizedSliceCover.lean`. Assembled short-parent and width hypotheses for synchronized slice covers. [Pinned source at the inspected commit](#).
- [180] [source 39] Same repository and snapshot, `Proofs16ClosingComparison.lean`. The numerical closing-comparison obstruction, not a counterexample to the full assertion of Lemma 16.10. [Pinned source at the inspected commit](#). *[write]* Source 39’s seven `Proofs16` files have the same blobs at HEAD `bec460a33` as at its pin (checked at the write); none of its research reports is cited, although source 19 had been placed (note in Section F.44).

Added in the eighth write (sources 40–46, 7 October 2026). Works above that sources 40–46 also cite: [1] (all seven; source 40 the DOI, “Section 15, especially Corollary 15.2 and Lemmas 15.4–15.5”; source 41 a copy under `~zeilberg`, “Section 5, especially printed pp. 489–493”; source 42 the same copy, Section 5; source 43 the DOI, “Section 17, especially Lemma 17.1 and Proposition 17.2; original printed page 578 inspected”; source 44 the DOI and the copy under `users/gasarch`, “Section 15, especially pp. 562–565”; source 45 the DOI and the copy under `~gasarch`; source 46 the DOI); [2] (sources 40 and 46; source 40 as version 2, 29 February 2024, “used only to distinguish the local refinement”); [72] (source 41, as arXiv:2011.12275v1, 24 November 2020, “Theorem 1.1, Corollary 1.2, and Lemma 6.1”); [101] (source 42, “Lemma 2 and Section 5”); [150] (sources 43 and 46; source 43 “Definition 3.1, Proposition 3.5, Lemma 4.1, and Proposition 4.6”); [98] and [99] (source 44, as background); [7] (source 45, as *Journal d’Analyse Mathématique* **117** (2012), 133–186). The entries below are new.

- [181] R. Bhatia and C. Davis, *A better bound on the variance*, *American Mathematical Monthly* **107** (2000), 353–357. [doi:10.1080/00029890.2000.12005203](https://doi.org/10.1080/00029890.2000.12005203) (source 42).
- [182] M. N. Kolountzakis and Sz. Gy. Révész, *Turán’s extremal problem for positive definite functions on groups*, arXiv:math/0312218, first version 10 December 2003; author PDF, October 2004 version, Theorem 2. <https://eigen-space.org/ps/turan-groups.pdf> (source 42).

- [183] J. Fox, M. W. Xu and Y. Zhou, *Discrepancy in modular arithmetic progressions*, *Compositio Mathematica* **158** (2022), 2082–2108; preprint 2021, <https://arxiv.org/abs/2104.03929> (source 42).
- [184] T. Tao and T. Ziegler, *The inverse conjecture for the Gowers norm over finite fields in low characteristic*, *Annals of Combinatorics* **16** (2012), 121–188. DOI: [10.1007/s00026-011-0124-3](https://doi.org/10.1007/s00026-011-0124-3); [arXiv:1101.1469](https://arxiv.org/abs/1101.1469) (sources 43 and 46). Not the Tao–Ziegler paper cited above as [77].
- [185] A. Prasad and M. K. Vemuri, *Decomposition of phase space and classification of Heisenberg groups*, [arXiv:0806.4064](https://arxiv.org/abs/0806.4064) (2008). Cited for the established finite-Heisenberg and symplectic-decomposition background (source 43).
- [186] A. Berger, A. Sah, M. Sawhney and J. Tidor, *Non-classical polynomials and the inverse theorem*, *Mathematical Proceedings of the Cambridge Philosophical Society* **173** (2022), 525–537. [doi:10.1017/S0305004121000682](https://doi.org/10.1017/S0305004121000682); [arXiv:2107.07495](https://arxiv.org/abs/2107.07495) (source 46).
- [187] L. Milićević, *General inverse theory for the U^4 norm*, preprint (2026), [arXiv:2601.01682](https://arxiv.org/abs/2601.01682) (source 46, “not exhaustively compared”). A different preprint from [32].
- [188] H. Kamiya, A. Takemura and H. Terao, *The characteristic quasi-polynomials of the arrangements of root systems and mid-hyperplane arrangements*, in *Arrangements, Local Systems and Singularities*, *Progress in Mathematics* **283**, Birkhäuser, 2010, pp. 177–190. <https://arxiv.org/abs/0707.1381>. Related modular-arrangement context (source 44).
- [189] N. Hegyvári, *Some remarks on the distribution of additive energy*, *Results in Mathematics* **78** (2023), Article 233. DOI: [10.1007/s00025-023-02010-5](https://doi.org/10.1007/s00025-023-02010-5); <https://arxiv.org/abs/2212.07109> (source 45).
- [190] E. Nazarewicz, M. O’Brien, M. O’Neill and C. Staples, *Equality in Pollard’s theorem on set addition of congruence classes*, *Acta Arithmetica* **127**(1) (2007), 1–15. Primary text: <https://www.impan.pl/shop/en/publication/transaction/download/product/82135> (source 45).
- [191] X. Shao, *Large values of the additive energy in $\mathbb{R}^d$ and $\mathbb{Z}^d$* , *Mathematical Proceedings of the Cambridge Philosophical Society* **156** (2014), 327–341. DOI: [10.1017/S0305004113000741](https://doi.org/10.1017/S0305004113000741); <https://arxiv.org/abs/1308.2247> (source 45).
- [192] [source 40] V. Reshetnikov and contributors, *ProveIt repository, Combinatorics/Ramsey*, materials inspected 6 October 2026, <https://github.com/VladimirReshetnikov/ProveIt/tree/main/Combinatorics/Ramsey>: the edited source `Papers/sz-thm-gowers-proof/sz-thm-gowers-proof.tex`, Section 15, and the index `Research/GowersSzemerédi/local-quantitative-refinements/README.md`. “Repository contents are mutable; this citation records the paths and inspection date rather than claiming a commit-pinned audit.”
- [193] [source 41] V. Reshetnikov, *ProveIt*, GowersSzemerédi source, commit `c94998a5db18ba63c815a3bf6c3877a0fde91930`: [Corollary 5.8 predicate](#); [coefficient polynomial definition](#).
- [194] [source 42] ProveIt research article, *Local quantitative refinements of Gowers’s proof of Szemerédi’s theorem* (this report), commit `b16ce380d7b82557b7120ca56c0a78391fdaf35e`, source-10 variance-to-density lemma, lines 4804–4846. “The exact source URL and bounded screening scope are in `SOURCES.md`” (shipped as `42-quadratic-scale-SOURCES.md`).
- [195] [source 42] Report275, “preceding research report on threshold-free affine density increments”, the theorem *Threshold free cube root bound* and its proof: source 26 of this report (Theorem I.41.1); identified in `42-quadratic-scale-SOURCES.md` by the SHA-256 of its manuscript.
- [196] [source 42] Report277, *Threshold free polynomial density increments on cyclic groups*, 6 October 2026, Section 5 and the retained affine bounds: source 41 of this report (Section I.62); identified by the SHA-256 of its manuscript.

- [197] [source 43] *Sharp Symmetry Certification for Gowers Derivative Spectra: Alternating-rank bounds, lossless polarization, and robust finite-field inverse steps*, research manuscript prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt project, 6 October 2026; “User Library file `Sharp_Symmetry_Certification.tex`”, question labels `ssc:q:two`, `ssc:q:nonclassical`, `ssc:q:composite`, recorded pin `6c2e2172b`: source 36 of this report (Section IV.15).
- [198] [source 43] V. Reshetnikov, *ProveIt*, Ramsey subtree, comparison snapshot `797aa3cfe967349322b806692a500f677f7047cf`, [pinned Ramsey directory](#): `Lean/GowersSzemerédi/Sections17_18.lean` and `Definitions.lean` at that revision; this directory’s README at blob `fe4d31eb3077580b2509ea146315d9857eb021d3` (shipped record: `data/43-boolean-phase-source_audit.json`).
- [199] [source 44] “Baseline A”: *The Codimension-One Degeneracies of Gowers Arrangements: An asymptotically sharp replacement for Lemma 15.4 and explicit random-selection thresholds*, AI-assisted research manuscript prepared for Vladimir Reshetnikov, 6 October 2026, inspected from the user’s research library as `arrangement_degeneracy.tex`: source 25 of this report (Section F.30).
- [200] [source 44] “Baseline B”: *Sharp Degeneracy Bounds for Gowers Arrangements: Hyperplane obstructions, second-order structure, and an exact finite-field count*, AI-assisted research draft prepared for Vladimir Reshetnikov and ProveIt, 6 October 2026, inspected library entries `article(20261006-231449).tex` and `article(20261006-231446).pdf`: source 27 of this report (Section F.32).
- [201] [source 44] V. Reshetnikov, *ProveIt: Combinatorics/Ramsey*, inspected snapshot `8c12c906e4ab5a00001a65de3c882f935a15661a`, <https://github.com/VladimirReshetnikov/ProveIt/tree/8c12c906e4ab5a00001a65de3c882f935a15661a/Combinatorics/Ramsey>; interface `Lean/GowersSzemerédi/Sections14_15.lean`, blob `5d91dce37bbca59fc76dfd30ec3a485d72ba583c`.
- [202] [source 45] ProveIt repository, *GowersSzemerédi: source transcription, formalization, and local quantitative refinements*, *Combinatorics/Ramsey*, inspected at commit `6ce33539ea468f51719a211c18aabdde262a24dd`, 7 October 2026 (UTC), <https://github.com/VladimirReshetnikov/ProveIt/tree/6ce33539ea468f51719a211c18aabdde262a24dd/Combinatorics/Ramsey>; report labels and source boundaries in its Appendix V.169 and `data/45-rigidity-gaps-source_manifest.json`. Printed from its BibTeX database.
- [203] [source 46] V. Reshetnikov, ProveIt repository, corrected Gowers-paper transcription, *Combinatorics/Ramsey/Papers/sz-thm-gowers-proof/*, [repository directory](#); inspected blob `41085a0efde8c1932e86e80c791984f52841c22e`.
- [204] [source 46] V. Reshetnikov, ProveIt repository, local quantitative refinements inventory, *Combinatorics/Ramsey/Research/GowersSzemerédi/local-quantitative-refinements/README.md*, [repository directory](#); inspected blob `fe4d31eb3077580b2509ea146315d9857eb021d3`. *[write]* The seven sources’ cited Lean and paper files have the same blobs at HEAD `6d4942222` as at their pins or recorded blobs (checked at the write).

Added in the ninth write (sources 47–53, 7 October 2026). Works above that sources 47–53 also cite: [1] (sources 47–50; source 47 “Section 17, especially Lemma 17.1 and Proposition 17.2”, with the copy under `~gasarch`; source 48 “Lemma 9.3 and its conditional sampling argument, printed pages 516–518”; source 49 the DOI; source 50 “the published pages 574–575, containing Lemma 16.10 and its proof”, with the copy under `users/gasarch`); [150] (source 47, as *Discrete Analysis* **2022**:14, “Definition 3.1 and Proposition 3.5”, background); [2] (source 48, version 2, “to distinguish the local constructive problem”); [98] (source 49, the finite-field arrangement method credited in its Section 5); [71] (source 49, as a preprint of 8 September 2026). Sources 51–53 have no bibliography; they cite the repository through pinned links. The entries below are new; [208] is not [34].

- [205] *Sharp Boolean Phase Integration from Gowers Derivative Spectra: Exact cubic obstruction ranks, repeated-variable certification, finite-abelian orthogonality, and a quartic reduction*, research manuscript

- prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt project, 6 October 2026; inspected library source `boolean_phase_integration.tex`, its comparison pin 797aa3cfe967. This is source 43 of this report (Section F.48) (source 47).
- [206] V. Reshetnikov, repository maintainer, *ProveIt: Combinatorics/Ramsey*, inspected repository interfaces and research materials, October 2026; observed branch ref 5c9a442d263632fdcd4e9690154c12c2dcc70b53. <https://github.com/VladimirReshetnikov/ProveIt/tree/main/Combinatorics/Ramsey> (source 47; not cited in its text).
 - [207] *Polynomial-loss restriction in Gowers’s Szemerédi argument: Fejér filters, density-sensitive moments, and finite-field systems*, research note prepared for the ProveIt project, 6 October 2026; inspected prior research file `gowers_polynomial_restriction.tex`, source audit anchored to dca9f06383ccdfc2ed208ccb7c9882214fdb09ac. This is source 17 of this report (Section F.22) (source 48).
 - [208] P. Raghavan, *Probabilistic construction of deterministic algorithms: approximating packing integer programs*, Journal of Computer and System Sciences **37** (1988), 130–143. doi:10.1016/0022-0000(88)90003-7. Cited for classical derandomization context (source 48).
 - [209] V. Reshetnikov, *ProveIt*, Ramsey development, snapshot 0b4890a1cd1b5417bcf0724db27e35b77bfa96bc; file `Combinatorics/Ramsey/Lean/GowersSzemerédi/Proofs09Restriction.lean` (source 48).
 - [210] V. Reshetnikov, *ProveIt: Combinatorics/Ramsey*, corrected Gowers transcription and project development, snapshot 8c12c906e4ab5a00001a65de3c882f935a15661a, inspected 6–7 October 2026 (source 49).
 - [211] *Local Quantitative Refinements of Gowers’s Proof of Szemerédi’s Theorem*, integrated sources 01–14, 6 October 2026, in the pinned research directory: this report at its third write (b16ce380d) (source 49).
 - [212] Research report prepared for Vladimir Reshetnikov with ChatGPT, *Degree Filtrations in Gowers’ Ramsey Framework*, 6 October 2026, `gowers_degree_stratification.tex`: source 22 of this report (source 49).
 - [213] Research report prepared for Vladimir Reshetnikov with ChatGPT, *Prime Torsion and Norm Gaps in Gowers Cube Counts*, 6 October 2026, `gowers_cube_structure.tex`: source 34 of this report (source 49).
 - [214] Research report prepared for Vladimir Reshetnikov with ChatGPT, *Graded Fourier Selection and Two Further Refinements of Gowers’ Argument*, 6 October 2026, `gowers_graded_refinements.tex`: source 35 of this report (source 49).
 - [215] Research report prepared for Vladimir Reshetnikov with ChatGPT, *Efficient Interpolation in Gowers’ Szemerédi Framework*, 6 October 2026, `gowers_interpolation.tex`: source 19 of this report (source 49).
 - [216] A. Björner and T. Ekedahl, *Subspace arrangements over finite fields: cohomological and enumerative properties*, preprint (1996), [arXiv:math/9612217](https://arxiv.org/abs/math/9612217) (source 49).
 - [217] W. Hoeffding, *Probability inequalities for sums of bounded random variables*, Journal of the American Statistical Association **58** (1963), 13–30. doi:10.1080/01621459.1963.10500830 (source 50).
 - [218] VladimirReshetnikov/ProveIt, *Combinatorics/Ramsey/Lean/GowersSzemerédi/Section16.lean*, commit 588552b79497c93d0e2f3a3f9224fd7681e8d396, blob c5c7d2bee91bc589de1d3082429f7ee3597e99f9 (source 50).

- [219] VladimirReshetnikov/ProveIt, `Definitions.lean`, same pinned commit; blob 97113b7afa6925a2dd4b76641eeaeff09597ab6a; definitions `IsMultilinear`, `Box`, `Box.width` (source 50).
- [220] VladimirReshetnikov/ProveIt, `Proofs16ClosingComparison.lean`, same pinned commit (source 50).
- [221] VladimirReshetnikov/ProveIt, `Sections14_15.lean`, same pinned commit; blob 5d91dce37bbca59fc76dfd30ec3a485d72ba583c; `HasProductProperty` and `AxisCube` (source 50).
- [222] VladimirReshetnikov/ProveIt, `Proofs16LocalAffineLift.lean`, same pinned commit; blob 03f86c7c8145fe3a1faa13953abb9f0c0bf6cb9e; the theorem `section16_local_affine_lift` (source 50).
- Added in the tenth write (sources 54–59, 7 October 2026).* Works above that sources 54–59 also cite: [1] (all six; source 54 “Proposition 8.1, journal page 511 (PDF page 47)”, source 58 “especially Lemma 9.3 and Corollary 9.4”, with the copies under `gasarch`); [2] (sources 54, 57, 58 and 59, as a global benchmark, not a dependency); [32] (source 54, “abstract and version metadata inspected”); [7] (source 55 in its journal version, *J. Anal. Math.* **117** (2012), 133–186, and source 57, version 2, “Theorem 1.1 and Remark 1.6”); [71] (source 55); [49] (sources 56 and 57); [189] (source 57); [4] (source 59, “published Lemmas 8.3 and 9.2; displayed probability estimate on p. 106”); [36], [12], [27] and [28] (source 59; it cites the Monthly version of [28], **121** (2014), 674–688). The entries below are new; [237] is not [129].
- [223] V. Reshetnikov, maintainer, *ProveIt: Combinatorics/Ramsey*, repository snapshot 6a6d7961c3c261da7dc18def13fff30b12dd5c04; the edited Gowers transcription, the Section 8 statement module, the research index and a scoped source audit are identified in its `SOURCE_AUDIT.md`, shipped as `54-single-phase-SOURCE_AUDIT.md` (source 54).
- [224] T. S. Motzkin and E. G. Straus, *Maxima for graphs and a new proof of a theorem of Turán*, Canadian Journal of Mathematics **17** (1965), 533–540. doi:10.4153/CJM-1965-053-6 (source 54).
- [225] Q. Tang, X. Zhang, C. Zhao and P. Zhao, *On the maxima of Motzkin–Straus programs and cliques of graphs*, Journal of Global Optimization **84** (2022), 989–1003. doi:10.1007/s10898-022-01187-3; the publisher abstract was inspected for its warning about erroneous maximizer classifications (source 54).
- [226] Z. Bao, J. Briët, D. Castro-Silva, P. van Dordrecht and J. Helsen, *On Clifford hierarchy testing and near-extremizers of noncommutative uniformity norms*, arXiv:2605.26983, version 1 (2026); consulted for the noncommutative extension question, not a proof input (source 55).
- [227] Vladimir Reshetnikov’s ProveIt repository, *Combinatorics/Ramsey*, especially *Research/GowersSzemerédi/local-quantitative-refinements*, inspected snapshot 128f514afce0bf7eb48133fb30dd3f8e034ea3a2, October 2026; “its research placement confers no proof-assistant status” (source 55).
- [228] D. Kazhdan and T. Ziegler, *Polynomial functions as splines*, Selecta Mathematica (N.S.) **25** (2019), Paper No. 31; preprint version inspected: arXiv:1712.09047v3, 4 December 2018, Theorems 1.6 and 1.7. <https://arxiv.org/abs/1712.09047> (source 56).
- [229] S. Peluse, *Another proof of the $U^4(\mathbb{F}_p)$ -inverse theorem*, arXiv:2609.15788v1, 14 September 2026, Section 7 (source 56).
- [230] V. Reshetnikov, *ProveIt*, <https://github.com/VladimirReshetnikov/ProveIt>, snapshot b6a8ee96631fa845067e9ec9e5b1c2f42960a9c9; the Section 6–7 and Section 10 statements, `Proofs10Shift.lean` and the research README, as recorded in its `source_manifest.json`, shipped as `data/56-relative-affine-repair-source_manifest.json` (source 56).
- [231] S. Huczynska, *Beyond sum-free sets in the natural numbers*, Electronic Journal of Combinatorics **21**(1) (2014), P1.21. doi:10.37236/2810 (source 57).

- [232] V. Reshetnikov, *ProveIt: Combinatorics/Ramsey*, repository snapshot 128f514afce0bf7eb48133fb30dd3f8e034ea3a2 (7 October 2026 UTC); the edited Gowers transcription, `FORMALIZATION_STATUS.txt`, `Lean/GowersSzemerédi/Sections12_13.lean` and the Section 13 extraction modules (source 57).
- [233] ProveIt research archive, source 40, *One-stage Fejér selection for Gowers arrangements: single-monomial relation isolation, polynomial parameter loss, and sharp exponents within a positive-Fourier selection class* (6 October 2026), archive docs/incoming/fejér-arrangement-selection.zip at commit 8ed496b14e46aad80a50fa7c34ea50d3c66a9323, blob 44036f557a4d017e4c198d35830c72ef3190e061. This is source 40 of this report (Section F.45) (source 57).
- [234] ProveIt research archive, source 45, *Rigidity gaps in Gowers’s Szemerédi framework: exact additive-energy levels, punctured cosets, and a real fourth-uniformity inequality* (7 October 2026), archive docs/incoming/Gowers_Rigidity_Gaps_Research.zip at commit 8ed496b14e46aad80a50fa7c34ea50d3c66a9323. This is source 45 of this report (Section F.50) (source 57).
- [235] I. D. Shkredov, *On Gowers norms of some functions*, arXiv:0912.2436v1, 12 December 2009, especially Lemma 2.1. <https://arxiv.org/abs/0912.2436> (source 58).
- [236] VladimirReshetnikov/ProveIt, *Local Quantitative Refinements of Gowers’s Proof of Szemerédi’s Theorem*, repository research report and companion materials, sources 04, 17, 29 and 40, inspected at commit 5d4718772d1a7dd4c48e1f49c8f92390c73b7ff1 (7 October 2026); this report and its directory, with the Lean interfaces under `Combinatorics/Ramsey/Lean/GowersSzemerédi` (source 58).
- [237] W. T. Gowers and L. Milićević, *A quantitative inverse theorem for the U^4 norm over finite fields*, 2017. <https://arxiv.org/abs/1712.00241>; the finite-field graph restriction used for comparison is Lemma 3.3 (source 59).
- [238] B. Bedert, *Riesz products and the Lonely Runner Conjecture: A wider gap of loneliness*, 2025 preprint. <https://arxiv.org/abs/2511.16636>; Theorem 1.3 and Lemma 4.1 in the version inspected on 7 October 2026 (source 59).
- [239] G. Perarnau and O. Serra, *The Lonely Runner Conjecture turns 60*, 2024, revised 2025. <https://arxiv.org/abs/2409.20160> (source 59).
- [240] V. Reshetnikov, ProveIt repository, *Combinatorics/Ramsey*, inspected at commit bec460a3335f8da421459b4bf69d082e674560a6; the combined report (sources 01–33 at that commit) and the README’s later-source audit, with blob identifiers in its `source_manifest.json`, shipped as `data/59-character-affine-source_manifest.json`. This is this report at its sixth write (source 59).
- Added in the eleventh write (sources 60–69, 7 October 2026).* Works above that sources 60–69 also cite: [1] (all ten; source 63 with the copy under `gasarch` and the corrected transcription, source 67 the same copy); [2] (sources 60, 67 and 69, as a global benchmark); [49] (source 60); [7] (sources 60, 63 and 67; source 60 in its journal version, *J. Anal. Math.* **117** (2012), 133–186, sources 63 and 67 version 2); [36] and [11] (source 61); [37] and [33] (source 64, the journal version of [37] and version 2 of [33], “28 July 2026”; source 67 the same version of [33]); [150] (sources 65 and 69); [190] (source 66); [152] (source 67, version 3). The entries below are new; [242] is cited by sources 60 (Theorem 1.2, equation (1.5)), 63 and 68 (Theorem 1.1, equation (1.8)).
- [241] B. Szegedy, *Limits of functions on groups*, arXiv:1502.07861 (2015). <https://arxiv.org/abs/1502.07861> (source 60).
- [242] M. Christ and M. Iliopoulou, *Inequalities of Riesz–Sobolev type for compact connected Abelian groups*, arXiv:1808.08368v2 (2019). <https://arxiv.org/abs/1808.08368v2> (source 60: Theorem 1.2,

- equation (1.5); source 63: compact-group context; source 68: Theorem 1.1 and equation (1.8), “the circle case is attributed” to Riesz–Sobolev).
- [243] Vladimir Reshetnikov’s ProveIt project, *Gowers–Szemerédi local quantitative refinements*, repository snapshot 89c1cbc38600279fed03c8e741aba9eafb08ca70, 7 October 2026, 04:12:28 UTC; the inspected README inventories sources 01–46 and the consolidated article contains sources 01–33 (source 60).
 - [244] ProveIt research collection, source 04, *Sharper Local Bounds in Gowers’s Proof of Szemerédi’s Theorem: polynomial homomorphism restriction, sharp cube estimates, and floor-pattern complexity*, consolidated in the snapshot [243]; original manuscript `gowers_quantitative_refinements.tex`. This is source 04 of this report (Section F.9) (source 60).
 - [245] ProveIt research collection, source 17, *Polynomial-loss restriction in Gowers’s Szemerédi argument: Fejér filters, density-sensitive moments, and finite-field systems*, consolidated in the snapshot [243]; original archive `Gowers_Polynomial_Restriction_Research.zip`. This is source 17 of this report (Section F.22) (source 60).
 - [246] ProveIt research collection, source 19, *Efficient Interpolation in Gowers’ Szemerédi Framework: exact sampling losses, shared refinements, and simultaneous polynomial recurrence*, consolidated in the snapshot [243]; original manuscript `gowers_interpolation.tex`. This is source 19 of this report (Section F.24) (source 60).
 - [247] ProveIt research collection, source 29, *Further Quantitative Refinements of Gowers’s Szemerédi Argument: higher norm comparisons, joint quadratic rigidity, and restriction by affine rank*, consolidated in the snapshot [243]; labels `gsr:cc:29:gmsr:moment-theorem`, `gsr:cc:29:gmsr:cd-upper`, `gsr:cc:29:gmsr:one-third`. This is source 29 of this report (Section F.34) (source 60).
 - [248] ProveIt research collection, source 33, *Sharp High-Agreement Stability for Gowers’s Respected Arrangements: full-product decoding, optimal error terms, and erasures*, 6 October 2026, consolidated in the snapshot [243]; its local theorem `gsr:fm:33:ras:thm:local-main`. This is source 33 of this report (Section F.38) (source 60).
 - [249] B. Green and I. Z. Ruzsa, *Freiman’s theorem in an arbitrary abelian group*, Journal of the London Mathematical Society (2007). doi:10.1112/jlms/jdl021; <https://arxiv.org/abs/math/0505198>; background, not a dependency of the principal proofs (source 61).
 - [250] PROVEIT repository, *GowersSzemerédi: Sections06_07.lean and Sections08_09.lean*, commit 5ea1d06dd2f2d3a30e96016b236e3f1e3dfb048f; blob identifiers beginning 57b79242466c and 0fb6e0a1de3f (source 61).
 - [251] PROVEIT repository, *GowersSzemerédi/Proofs09Restriction.lean*, same commit; blob identifier beginning 29bb4178fbec; contains `lemma_9_3_holds`, “inspected, not independently compiled” (source 61).
 - [252] PROVEIT repository, *Local quantitative refinements of Gowers’s proof of Szemerédi’s theorem: research README*, same commit; used for scope and overlap comparison. This is this directory’s README at its sixth write (source 61).
 - [253] J. B. Rosser and L. Schoenfeld, *Approximate formulas for some functions of prime numbers*, Illinois Journal of Mathematics **6** (1962), 64–94, Theorem 15. doi:10.1215/ijm/1255631807 (source 62).
 - [254] V. Reshetnikov, repository maintainer, *ProveIt: Local Quantitative Refinements of Gowers’s Proof of Szemerédi’s Theorem*, research collection README at revision 17123532a2de638477dce1c92eb0b0b20a247d2f, inspected 6 October 2026; “the source-17 and source-29 summaries distinguish prime/fixed-field restriction from the all-moduli interface” (source 62).

- [255] ProveIt contributors, *Sections08_09.lean*, GowersSzemerédi development, revision 17123532a2de638477dce1c92eb0b0b20a247d2f (source 62).
- [256] ProveIt contributors, *Proofs09Restriction.lean*, GowersSzemerédi development, same revision (source 62).
- [257] T. Tao, *Higher Order Fourier Analysis*, Graduate Studies in Mathematics 142, American Mathematical Society, 2012; author’s preliminary text at <https://terrytao.wordpress.com/wp-content/uploads/2012/12/gsm-142-tao7-higher-book-05june2012.pdf> (source 63).
- [258] V. Reshetnikov, repository maintainer, *ProveIt: Combinatorics/Ramsey*, research collection and corrected Gowers transcription, snapshot d7b83fa1fe60478e3ed8525d58f747a1d76e35c2, inspected 7 October 2026 (source 63).
- [259] ChatGPT, prepared for Vladimir Reshetnikov, *Local Phase Removal and Cubic Rigidity in Gowers’ Szemerédi Framework*, unpublished research manuscript, 7 October 2026; “inspected predecessor TeX file, 126,709 bytes”, SHA-256 9cd0a536250370bb2126de79bc063ba20d6c48e351b5224f08697b95f3a24343; the source of the sliding-window positive-cover interface and the denominator $\widehat{D}_k$ (source 63). *[write, 7 October 2026; sources 60–69]* This manuscript was never delivered to this repository; it is not source 16, *Local Phase Removal Without Carry Patterns* (Section F.4).
- [260] W. T. Gowers, *A uniform set with fewer than expected arithmetic progressions of length 4*, Acta Mathematica Hungarica **161** (2020), 756–767. doi:10.1007/s10474-020-01072-z; arXiv:2004.07598 (source 64).
- [261] M. B. Nathanson, *Sidon sets for linear forms*, Journal of Number Theory **239** (2022), 207–227. doi:10.1016/j.jnt.2021.08.005; arXiv:2101.01034 (source 64).
- [262] Y. C. Cheng, *Greedy Sidon sets for linear forms*, Journal of Number Theory **266** (2025), 225–248. doi:10.1016/j.jnt.2024.07.010; background, not an input to the proofs (source 64).
- [263] *Sharp Boolean Phase Integration from Gowers Derivative Spectra: Exact cubic obstruction ranks, repeated-variable certification, finite-abelian orthogonality, and a quartic reduction*, research manuscript prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt project, 6 October 2026; inspected as boolean_phase_integration.tex with 43-boolean-phase-FORMALIZATION.md; comparison snapshot 797aa3cfe967349322b806692a500f677f7047cf. This is source 43 of this report (Section F.48) (source 65).
- [264] *Exact Boolean Obstruction Energies: Derivative periods, sharp quartic–sextic integration, and dimension-free higher-degree extremizers*, research manuscript prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt project, 6 October 2026, 23 pages; inspected as article(20261007-014512).tex; its repository snapshot 5c9a442d263632fdcd4e9690154c12c2dcc70b53. This is source 47 of this report (Section F.52) (source 65).
- [265] V. Reshetnikov, repository maintainer, *ProveIt: Combinatorics/Ramsey*, inspected interfaces at snapshot 8c40adc24df2e6f338cea728887772aead4bf144 (source 65; listed in its bibliography and not cited in its text).
- [266] ProveIt, source 45, *Internal AI proof cross-check: ordered-energy results*, file 45-rigidity-gaps-review_energy.md of this directory, blob 7984f49e51f4cfb0e9cd44f8385b5302bc75242f; “not an external referee report or a formal verification” (source 66).
- [267] V. Reshetnikov, maintainer, *ProveIt: Combinatorics/Ramsey*, inspected 7 October 2026; the merged research report at blob f4e5705da5d349256757ee7063c6a4480638d0d2 (source 67).

- [268] ProveIt research series, source 37, *Sharp Spectral Detection of Missing Intervals: exact moment tradeoffs, extremal rigidity, and quantitative recurrence in Gowers’s argument*, in the merged report; original `sharp_missing_intervals.tex`. This is source 37 of this report (Section F.42) (source 67).
- [269] ProveIt research series, source 39, *Optimal Shared-Anchor Interpolation and Labelled Slice Refinement in Gowers’s Szemerédi Argument: exact finite minimax bounds, sharp Poisson limits, and a replacement for the sampling step of Lemma 16.10*, in the merged report. This is source 39 of this report (Section F.44) (source 67).
- [270] ProveIt research series, source 11, *Sharp Affine Rigidity from Graph Additive Energy: cubic stability, target torsion, and exact non-affine extremizers*, in the merged report. This is source 11 of this report (Section F.16) (source 67).
- [271] ProveIt project working manuscript, *Sharp Character Separation and Affine Rigidity in Gowers’ Inverse Steps: torsion-sensitive Freiman restrictions, signed energy inequalities, and a research programme*, prepared with ChatGPT for Vladimir Reshetnikov, 7 October 2026; `gowers-character-affine-refinements.tex`, “inspected separately from the merged sources 1–39”; credited for the ranges $4/9$ and $1/2$, the recovery range $2/9$ and the cutoff $N \geq 16$. This is source 59 of this report (Section F.64) (source 67).
- [272] ProveIt research series, sources 10, 14, 22 and 38, polynomial degree filtrations, norm colourings and six-term norm defects, in the merged report; in particular source 22, *Degree Filtrations in Gowers’ Ramsey Framework*, and source 38, *Exact Quotient Reduction and Sharp Stability in Gowers’ Ramsey Framework* (source 67).
- [273] T. Mittal and S. Roy, *A General Framework for Low Soundness Homomorphism Testing*, arXiv:2509.05871, version 1, 2025. <https://arxiv.org/html/2509.05871v1> (source 67).
- [274] A. Weil, *On some exponential sums*, Proceedings of the National Academy of Sciences USA **34** (1948), 204–207. <https://doi.org/10.1073/pnas.34.5.204> (source 67).
- [275] R. Lidl and H. Niederreiter, *Finite Fields*, second edition, Encyclopedia of Mathematics and its Applications 20, Cambridge University Press, 1997, Theorem 5.41. <https://doi.org/10.1017/CB09780511525926> (source 67).
- [276] V. Reshetnikov, *ProveIt: Combinatorics/Ramsey*, repository snapshot 0d93c1eac2e20c313e7ed0cb7cb74ff705b00101, inspected 7 October 2026; source boundaries in its `data/source_manifest.json`, shipped as `data/68-sharp-energy-source_manifest.json` (source 68).
- [277] ProveIt research source 04, *Sharper Local Bounds in Gowers’s Proof of Szemerédi’s Theorem: polynomial homomorphism restriction, sharp cube estimates, and floor-pattern complexity*, in the consolidated manuscript of [276]; labels `gsr:en:04:rest:main`, `gsr:en:04:rest:eight`. Source 04 of this report (source 68).
- [278] ProveIt research source 14, *Sharp Extremal Refinements in Gowers’s Szemerédi Framework: secondary cube asymptotics, arrangement stability, and polynomial supersaturation*, in the consolidated manuscript of [276]; label `gsr:en:14:thm:energy-sharp-stability`. Source 14 of this report (source 68).
- [279] ProveIt research source 17, *Polynomial-loss restriction in Gowers’s Szemerédi argument: Fejér filters, density-sensitive moments, and finite-field systems*, in the consolidated manuscript of [276]; label `gsr:en:17:thm:main`. Source 17 of this report (source 68).
- [280] ProveIt research source 45, *Rigidity gaps in Gowers’s Szemerédi framework: exact additive-energy levels, punctured cosets, and a real fourth-uniformity inequality* (7 October 2026), full manuscript from the archive `docs/incoming/Gowers_Rigidity_Gaps_Research.zip` at commit 8ed496b14e46aad80a50fa7c34ea50d3c66a9323. Source 45 of this report (Section F.50) (source 68).

- [281] ProveIt research source 58, *Polynomial restriction on every cyclic group: a torsion-sensitive refinement of Gowers’s Lemma 9.3 and an optimal exponent for positive Fourier filtering* (7 October 2026), full manuscript from `docs/incoming/gowers_torsion_uniform_restriction.zip` at commit `f8bc5e2ec281fcee02c0fc5ef717a775ee7e4442`. Source 58 of this report (Section F.63) (source 68).
- [282] ProveIt research source 59, *Sharp Character Separation and Affine Rigidity in Gowers’ Inverse Steps: torsion-sensitive Freiman restrictions, signed energy inequalities, and a research programme for the ProveIt Ramsey project* (7 October 2026), full manuscript from `docs/incoming/gowers-character-affine-refinements.zip` at commit `f8bc5e2ec281fcee02c0fc5ef717a775ee7e4442`. Source 59 of this report (Section F.64) (source 68).
- [283] *Sharp Boolean Phase Integration from Gowers Derivative Spectra: Exact cubic obstruction ranks, repeated-variable certification, finite-abelian orthogonality, and a quartic reduction*, research manuscript prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt project, 6 October 2026; “Source 43 in the project collection”. Source 43 of this report (source 69).
- [284] *Exact Boolean Obstruction Energies: Derivative periods, sharp quartic–sextic integration, and dimension-free higher-degree extremizers*, research manuscript prepared with ChatGPT for Vladimir Reshetnikov and the ProveIt project, 6 October 2026; “Source 47 in the project collection”, inspected as `article(20261007-014512).tex`, its comparison revision `5c9a442d263632fdcd4e9690154c12c2dcc70b53`. Source 47 of this report (source 69).
- [285] V. Reshetnikov, repository maintainer, *ProveIt: Combinatorics/Ramsey*, research collection and formal interfaces; “repository hosting does not give this manuscript formal proof status” (source 69).