

Aligned Fragments and Constrained Crossover

Undecidability, exact generation depth,
and rational growth of context-free recombination closures

Research draft prepared with ChatGPT
for Vladimir Reshetnikov

29 September 2026

Abstract

For equal-length words, constrained one-point crossover exchanges suffixes at the same position. We prove that deciding whether a context-free language satisfies $L \otimes_c L = L$ is co-recursively-enumerable-complete, already over a fixed three-letter alphabet. This answers the decision question explicitly left open in Section 4 of Hughes, arXiv:2608.27755v1. A delimiter amplification also proves undecidability of eventual finite stabilization, with an exact rank- $r + 1$ witness built from r omitted blocks.

The structural invariant is the least number $\rho_L(w)$ of consecutive fragments of w that occur at the same positions in seed words of the same length. For parallel self-crossover we prove $\rho_{L \otimes_c L}(w) = \lceil \rho_L(w)/2 \rceil$ and hence exact generation depth $\lceil \log_2 \rho_L(w) \rceil$. Rank spectra have no gaps, and every length- n slice stabilizes by $\lceil \log_2 n \rceil$, sharply.

For regular inputs, one-step closure has a polynomial-time DFA test and is PSPACE-complete for NFAs; finite stabilization reduces to the classical limitedness problem for distance automata. For every context-free seed, the full closure has an effectively rational commutative generating function. Nevertheless, an explicit binary linear context-free seed has a closure that is not context-free. Its finite generations have polynomial growth and minimal recurrence order $q2^k + 1$, whereas the limit has exponential growth. Complete proofs, exact finite audits, and a research agenda are included.

Status. Unrefereed research draft. The proofs are mathematical arguments, not Lean or Rocq certificates. A bounded literature search supports the connection to the stated open question, not a guarantee of first discovery for every structural lemma or consequence.

[write] A report of the ProveIt research-report collection (category automata-and-formal-languages, report constrained-crossover-closure), written from batch-44 manuscript 02; text marked [write] was added then. The sibling report on insertion-degree spectra is related in Section 1.4; provenance is in Appendix D. [Added 30 September 2026, batch 66: Part II (from Section 13) proves finite stabilization of regular seeds PSPACE-complete for NFA input; Part III (from Section 29) decides regularity of the full closure of context-free seeds and classifies sparse one-marker seeds. Both come from manuscripts of batch 66; their opening sections give their provenance.]

Contents

Part I. Aligned fragments and constrained crossover	5
1 Problem, provenance, and principal results	5
1.1 A specific question behind the repository	5
1.2 What the theorems say	5
1.3 Attribution and limits of the novelty claim	6
1.4 [write] Relation to the sibling report on insertion-degree spectra	6
1.5 [write] Notation beside the neighbouring texts	7
2 Conventions and the coordinate hull	7
3 Aligned fragments and exact generation depth	9
3.1 The fragment invariant	9
3.2 There are no missing finite ranks	10
3.3 Sharp stabilization bounds	10
3.4 The frozen-source convention is also exactly solvable	11
4 One-step equality is undecidable for context-free inputs	11
4.1 A guard construction with universal first generation	12
5 Finite stabilization and an exact block amplification	13
6 Polynomial-time rank computation for a fixed target	14
7 Regular inputs: finite-state tests and limitedness	15
7.1 One-step closure for deterministic automata	15
7.2 Encoding all fragment ranks by a distance automaton	16
8 Rational counting for the full closure of every CFL	17
8.1 Effective semilinear position supports	17
8.2 A rational commutative generating function	18
9 A linear context-free seed with a non-context-free closure	20
9.1 Definition, grammar, and position geometry	20
9.2 Every rank, every generation, and exact counts	21
9.3 Sharp recurrence orders of the finite generations	22
10 Verification, reproducibility, and proof boundaries	23
10.1 Executed exact audits	23
10.2 Files and rebuilding	23
10.3 What was not verified or claimed	24
11 Further research questions and a formalization plan	25
11.1 A realistic Lean development	27
12 Conclusion	27
Part II. Finite stabilization of regular seeds is PSPACE-complete	29
13 [write, batch 66] Provenance and notation of Part II	29
14 The question and the contribution	31

15 Crossover, aligned fragments, and the two clocks	33
16 The exact-length geometry of an NFA	35
17 Uniform periodicity and its quantitative cost	36
18 Periodic interiors and the obstruction dichotomy	36
19 Polynomial-space decision and finite obstruction bounds	39
20 Binary PSPACE-hardness	41
21 Deterministic automata and safe cores	43
22 Quantitative consequences and worked examples	44
23 Executable audits and their limits	45
24 Further research questions	48
25 A formalization and implementation roadmap	50
26 Conclusion	51
27 Proof and edge-case audit	51
28 Source ledger and novelty boundary	52
 Part III. Decidable regularity and exact stabilization of crossover closures	 54
29 [write, batch 66] Provenance and notation of Part III	54
30 Research target and relation to ProveIt	56
31 Coordinate supports and the full closure	58
32 Closure comparison and Boolean length theory	59
33 Decidable regularity through row types	60
34 One-marker seeds: exact generations	62
35 Sparse semilinear masks and arithmetic rays	63
36 A finite-ray pumping principle	64
37 The sparse-marker language-class trichotomy	65
38 An exact compatibility-graph formula	67
39 Exact-depth complexity for succinct regular seeds	68
40 Explicit examples and arbitrary finite depths	69
41 Verification, implementation, and proof status	71
42 Further research questions	73

43 Conclusion	74
44 Proof audit and boundary cases	75
45 Package and reproducibility	76
 Part IV. Quadratic finite rank for constrained crossover	 78
46 [write, batch 71] Provenance and notation of Part IV	78
47 The question and the result	80
48 Aligned fragments and generation depth	81
49 A length gate with a quadratic last gap	82
50 Turning missing lengths into rank	83
51 A six state example	85
52 The inherited upper bound and its consequences	85
53 Attribution and remaining questions	86
54 Reproducible finite checks	87
 Part V. Rational rank slopes for constrained crossover	 89
55 [write, batch 73] Provenance and notation of Part V	89
56 The question and the setting	91
57 Why longest aligned fragments are optimal	92
58 Viable layers and a fixed word algorithm	93
59 The periodic interior	94
60 The finite weighted reset graph	95
61 Cycle means with unrestricted endpoints	95
62 Exact rational slopes on the original residues	96
63 Examples and interpretation	97
64 A sharper finite rank bound	98
65 Reproducibility and remaining questions	100
 Part VI. Finite crossover stabilization for deterministic automata	 102
66 [write, batch 73] Provenance and notation of Part VI	102
67 Results and precise conventions	105

68 Rank and the inherited periodic criteria	105
69 Compressing safe cores into SCC stripes	107
70 A polynomial certificate for nonstabilization	108
71 Binary DFA hardness via length testers	109
72 A deterministic family with quadratic finite rank	110
73 A sharp leading constant for every complete DFA	112
74 Finite verification and attribution	115
 Part VII. Sharp leading constant for constrained crossover rank	 117
75 [write, batch 73] Provenance and notation of Part VII	117
76 Aligned fragments and the main results	120
77 Robust compression inside a primitive component	121
78 Stable phases and a universal cyclic component	125
79 Transferring the component bound to the whole NFA	127
80 The matching leading coefficient from a length gate	130
81 The leading coefficient for finite alphabets	130
82 A sharp linear family and its slice maxima	135
83 An exact algorithm for the maximal rank	136
84 Computer assisted binary extrema through four states	139
85 Sharper bounds from common paths	141
86 Graphs with a cycle on all but one vertex	144
87 A large exponent without a total labeled state subset	146
88 Scope and reproducible checks	147
 Appendices to Part I	 150
A The classical universality input	150
B Why the positional counts are eventually affine	150
C A compact proof-dependency map	151
D [write] Provenance of this report	151

1 Problem, provenance, and principal results

1.1 A specific question behind the repository

The ProveIt report *Gaps in Bounded-Shuffle Hierarchies*, including its later binary-profile extension, studies insertion-degree spectra and cites Hughes’s recent work on language operations [2, 1]. The report distinguishes its spectrum results from the surrounding decision problems. We follow that connection, but study a different operation: *equal-length, aligned, one-point crossover*.

Hughes defines constrained crossover on page 2 and, on page 3, explicitly leaves the decidability of $L \otimes_c L = L$ open for context-free languages. The obstruction to his general theorem is precise: constrained crossover does not contain concatenation. Our reduction does not try to restore concatenation. Instead, it makes every omitted word obtainable by a single aligned repair using a fresh guard letter. The decision problem is thereby settled directly.

[write] *Where the question is posed.* Checked when this report was written (29 September 2026) against the 13-page PDF of arXiv:2608.27755v1, whose printed page numbers are its PDF page numbers. Hughes defines unconstrained crossover $\otimes_u$ and constrained crossover $\otimes_c$ (the same symbol $\otimes_c$ as here) in Section 2 (“Definitions and Notations”, p. 2). The question is posed in Section 4 (“Foundational Results”), under “Single step equality”. His Theorem 1 makes single-step equality $L \odot L = L$ at least as hard as universality for every operator $\odot$ with $L \cdot L \subseteq L \odot L$ (on a class of languages for which containment of every letter and of ε is decidable but universality is not), and its Corollary applies this to simple insertion, shuffle and unconstrained crossover. The note after the Corollary (p. 3) observes that constrained crossover fails this subsumption criterion, so “the complexity of the question, ‘does $L \otimes_c L = L$ ’, is open.” The list of open questions that follows repeats it as its second item, for a context-free L (top of p. 4). Theorem 4.2 answers that question: the problem is co-r.e.-complete (Π_1^0 -complete), hence undecidable, already over a fixed three-letter alphabet. The first item of Hughes’s list, whether the shuffle of two context-free languages is context-free, is not addressed here. Hughes’s Section 14 (p. 12) also lists “Study constrained crossover” among its open directions without posing a specific question; Section 11 records how the present results bear on his other Section 14 questions.

The inspected ProveIt revision is

9b24a3a8d545af9624f6ac455f5b548be62818b6.

The relevant source is

SetTheory/Cardinals/docs/reports/
automata-and-formal-languages/insertion-degree-spectra/article.tex.

At this revision the source has both the original spectrum paper and its binary-profile Part II. We do not reprove their arbitrary-spectrum realization as a claimed contribution, and we do not resolve their both-singleton insertion conjecture. No theorem in the present paper assumes the correctness of an unverified ProveIt mathematical claim. [write] That source is the sibling report *insertion-degree-spectra* of the same collection category; its directory did not change between the pin and the placement of this report (commit 203016015). Section 1.4 states how the two reports relate.

1.2 What the theorems say

The results have three interacting parts. The first is a precise resource calculus: each generation halves an intrinsic fragment rank. The second is an algorithmic boundary: global stabilization is undecidable for

grammars but decidable for finite automata. The third is an enumeration theorem: the commutative counting series of a context-free seed’s full recombination closure is rational even when the resulting language is not context-free.

Question	Result proved here	Section
Exact generation of a word	$L^{[k]} = \{w : \rho_L(w) \leq 2^k\}$; a no-gap rank spectrum	3
One-step equality for CFGs	Co-r.e.-complete over a fixed ternary alphabet	4
Eventual equality for CFGs	Undecidable over four letters; strictness witnesses at every stage	5
A fixed target and grammar	Exact rank and generation decidable in polynomial time	6
Regular-input decision problems	DFA polynomial one-step test; NFA PSPACE-completeness; decidable finite stabilization	7
Counting the full closure of a CFL	Effective rationality, including commuting letter weights	8
Closure need not be context-free	Explicit linear seed; exact counts and sharp recurrence orders	9

1.3 Attribution and limits of the novelty claim

The conjectural-status claim is deliberately narrow: the one-step decision question is explicitly open in the inspected Hughes version. The proofs below answer that question under exactly its equal-length, equal-cut definition. The coordinate-product closure is an elementary recombination fact; it is reproved for completeness and is not advertised as a first discovery. Nor are context-free universality, effective Parikh semilinearity, Presburger counting, or distance-automaton limitedness new here.

The proposed contribution consists of the guarded reduction, its exact amplification, the quantitative fragment calculus and algorithms, and the rational-growth consequences in this setting. Equivalent formulations in recombination, splicing, or genetic-algorithm literature have not been exhaustively excluded. The article gives proofs so that its mathematical content can be assessed independently of priority. The source search and verification boundary are recorded again in Section 10.

1.4 [write] Relation to the sibling report on insertion-degree spectra

The sibling report *Gaps in Bounded-Shuffle Hierarchies* (ProveIt path `SetTheory/Cardinals/docs/reports/automata-and-formal-languages/insertion-degree-spectra/`) and this report start from the same preprint of Hughes, but from different parts of it.

- The sibling’s Part I answers questions of Hughes’s Sections 10–11. It refutes the insertion-degree no-gap conjecture (every finite set containing 1 is a spectrum), answers the finite representation-erasure question with a binary example, and proves an iteration-depth no-gap theorem (its Theorem 8.2). Its Part II realizes every finite spectrum containing 1 over a binary alphabet with one inserted unary word. Its scope sentence, under its abstract, says that it addresses “the indicated no-gap questions, not the surrounding undecidability results”. It never names the crossover question: the word “crossover” occurs there only in the title of Hughes’s preprint.
- This report answers the single-step question for constrained crossover posed in Hughes’s Section 4 (Section 1 above), one of those surrounding decision problems, for an operation the sibling does not study.
- No theorem of either report is used by the other, and none is proved in both. The no-gap theorem for aligned rank (Theorem 3.4) is an analogue, for a different operation, of the sibling’s iteration-depth no-gap theorem; neither implies the other. The sibling iterates $X_{m+1} = A \sqcup X_m$ with the inserted language A

held fixed. Its closest crossover counterpart is therefore the frozen-source iteration of Theorem 3.8, not the parallel iteration of Definition 2.1.

- Nothing here bears on the questions the sibling leaves open, in particular Hughes’s singleton insertion-degree no-gap conjecture (his Conjecture 2, the sibling’s Question 21.1), as the paragraph after Theorem 3.4 already says.

The sibling report carries a dated remark, under its scope sentence, that points to this report.

1.5 [write] Notation beside the neighbouring texts

No symbol of the manuscript was renamed. The table lists the readings that Hughes’s preprint or the sibling report could suggest.

Here	Meaning here	Not to be confused with
$A \otimes_c B$	constrained one-point crossover (Definition 2.1); the same operation and symbol as Hughes’s $\otimes_c$	Hughes’s unconstrained crossover $\otimes_u$, which contains concatenation and is covered by his Theorem 1
$T(L), L^{[k]}$	parallel self-crossover: both parents from generation k	Y_k , with one parent frozen in L (Theorem 3.8). Hughes’s iterated insertion starts from B and applies A on the left each time, and his iterated crossovers are defined “as with other operations”; for $A = B = L$ that reading gives Y_k , not $L^{[k]}$
$\rho_L(w)$	aligned-fragment rank (Definition 3.1)	the sibling’s insertion degree $d_{A,B}(w)$, a number of inserted pieces; and the unspecified “rank condition” of Hughes’s Section 14
$d_L(w)$	generation depth $\lceil \log_2 \rho_L(w) \rceil$ (5)	the sibling’s $d_{A,B}(w)$ (a degree, not a depth) and its iteration depth (a number of applications of $A \sqcup \cdot$)
“no gaps”	the ranks on a slice form $\{1, \dots, R_L(n)\}$ (Theorem 3.4)	the sibling’s spectra $\mathcal{D}(A, B)$, which <i>can</i> have gaps, and its iteration depths, which cannot
stabilization	some pair of consecutive generations is equal (Corollary 3.5)	Hughes’s “adjacent equality”; by Lemma 2.2 (and likewise for Y_k) adjacent equality here is already his “fixed point”
co-r.e.	negative instances effectively enumerable; co-r.e.-complete means Π_1^0 -complete	r.e. (the positive instances)

2 Conventions and the coordinate hull

All alphabets are finite. Put $\mathbb{N} = \{0, 1, 2, \dots\}$ and write ε for the empty word. For a word w of length n and $0 \leq i \leq j \leq n$, $w[i : j]$ denotes positions $i + 1, \dots, j$. Thus the indexing of boundaries is zero-based, but the indexing of letters is one-based. Write $L_n = L \cap \Sigma^n$.

Definition 2.1 (Constrained crossover). For languages $A, B \subseteq \Sigma^*$, define

$$A \otimes_c B = \{wz, yx : wx \in A, yz \in B, |w| = |y|, |x| = |z|\}.$$

For self-crossover put $T(L) = L \otimes_c L$. Its *parallel iterates* are

$$L^{[0]} = L, \quad L^{[k+1]} = T(L^{[k]}).$$

Both parents in generation $k + 1$ may be chosen from generation k .

The two length equalities imply that the complete parents have the same length and that their cut positions agree. Cuts at either endpoint are allowed. On a fixed slice,

$$T(L)_n = \bigcup_{i=0}^n \text{Pref}_i(L_n) \text{ Suff}_{n-i}(L_n). \quad (1)$$

The prefix and suffix in this product may come from different parents, but both parents must lie in L_n .

Lemma 2.2. *The map T is monotone and inflationary. Consequently the iterates increase, and equality at one pair of consecutive generations implies equality at all subsequent generations.*

Proof. Every pair of parents available in a smaller language is available in a larger one. Also, crossing a word with itself returns it. Thus $L \subseteq T(L)$. If $L^{[k]} = L^{[k+1]}$, applying the same map T to both sides proves the next equality, and induction finishes the argument. $\square$

Remark 2.3 (Do not change the iteration rule). The iteration $X_{k+1} = L \otimes_c X_k$, with one parent frozen in the original language, is not the iteration in Definition 2.1. In particular its resource accounting need not double each generation. The self-operation above is inflationary even though an iteration involving two different input languages can lose words of lengths absent from one input. These distinctions remove any ambiguity suggested by general remarks about nonmonotone constrained crossover.

[write] The general remark meant is in Hughes's Section 2 (p. 2): the series produced by iterated constrained crossover, unlike that of unconstrained crossover, can drop earlier elements from later sets. With two different inputs, as in $X_{m+1} = A \otimes_c X_m$ with $X_0 = B \neq A$, it does happen: the words of B whose length no word of A has are lost at the first step. For self-crossover it cannot happen under either clock: the parallel iterates increase by Lemma 2.2, and the frozen-source iterates of Theorem 3.8 increase by its formula.

Definition 2.4 (Coordinate hull). For $n \geq 1$ and $1 \leq i \leq n$, let

$$P_{n,i}(L) = \{u_i : u \in L_n\}.$$

Define $\text{Rec}(L)_n = \emptyset$ if $L_n = \emptyset$, and otherwise define

$$\text{Rec}(L)_n = P_{n,1}(L)P_{n,2}(L) \cdots P_{n,n}(L).$$

At length zero put $\text{Rec}(L)_0 = L_0$. Finally $\text{Rec}(L) = \bigcup_{n \geq 0} \text{Rec}(L)_n$.

The explicit convention at length zero is necessary. The empty product of coordinate sets must not introduce ε when the seed does not contain it. Every crossover preserves membership in the coordinate hull, since it introduces no new letter at any position. The converse reachability statement will follow from a stronger exact-rank theorem.

3 Aligned fragments and exact generation depth

3.1 The fragment invariant

Definition 3.1 (Aligned fragment and rank). Fix $w \in \Sigma^n$, $n \geq 1$. A nonempty interval $[i, j]$ is L -compatible for w if there is $u \in L_n$ with $u[i : j] = w[i : j]$. The rank $\rho_L(w)$ is the least positive integer r for which there is a partition

$$0 = i_0 < i_1 < \dots < i_r = n \quad (2)$$

whose r intervals are L -compatible. If there is no such partition, put $\rho_L(w) = \infty$. Set $\rho_L(\varepsilon) = 1$ when $\varepsilon \in L$, and ∞ otherwise.

The seed witnessing one fragment need not witness another. However, every witness is a complete seed word of length n , and every fragment is at its *original position*. Ordinary factor occurrence without alignment would give a different and incorrect invariant.

We have $\rho_L(w) = 1$ exactly when $w \in L$. For nonempty w , finite rank is at most $|w|$, because one may split any compatible partition into single-letter intervals. The existence of such a singleton partition is exactly membership in the coordinate hull.

Theorem 3.2 (Exact rank contraction). *For every language L and every word w ,*

$$\rho_{T(L)}(w) = \left\lceil \frac{\rho_L(w)}{2} \right\rceil, \quad (3)$$

where the right side is ∞ when $\rho_L(w) = \infty$.

Proof. First suppose w is nonempty. A fragment matching a word of $T(L)$ crosses at most one cut of its two seed parents. Splitting the fragment at that cut, when the cut lies strictly inside it, gives at most two nonempty L -compatible fragments. A partition of w into s $T(L)$ -compatible fragments therefore yields an L -compatible partition of size at most $2s$. In particular finite $\rho_{T(L)}(w)$ implies finite $\rho_L(w)$ and

$$\rho_{T(L)}(w) \geq \lceil \rho_L(w)/2 \rceil.$$

Conversely, take a shortest partition (2) of size r . Group consecutive fragments in pairs. If two adjacent fragments are witnessed by $u, v \in L_n$, cross u with v at their common boundary. The resulting word agrees with w throughout their union. A leftover final fragment already matches a seed, and seeds belong to $T(L)$. Thus the grouped partition has $\lceil r/2 \rceil$ $T(L)$ -compatible pieces. This proves the reverse inequality. The same argument rules out a finite rank on either side when the other side is infinite. For $w = \varepsilon$, crossover neither creates nor removes the empty word, so the convention gives the stated identity. $\square$

Theorem 3.3 (Generation formula and full closure). *For every $k \geq 0$,*

$$L^{[k]} = \{w : \rho_L(w) \leq 2^k\}, \quad \bigcup_{k \geq 0} L^{[k]} = \text{Rec}(L). \quad (4)$$

A reachable word of rank r first occurs at generation

$$d_L(w) = \lceil \log_2 r \rceil. \quad (5)$$

Consequently $T(L) = L$ if and only if every nonempty slice L_n is its full Cartesian product of coordinate projections.

Proof. Iterating Theorem 3.2 gives $\rho_{L^{[k]}}(w) = \lceil \rho_L(w)/2^k \rceil$ for finite ranks. Membership in $L^{[k]}$ is equivalent to rank one relative to that language, which proves the first formula and the depth formula. Finite rank is equivalent to coordinate-hull membership, proving the union formula. The hull is fixed by crossover because each output letter stays in its designated coordinate set. If L is itself fixed, all its iterates equal L and therefore so does their union. The converse follows from the hull's fixedness. $\square$

The proof can also be viewed as a balanced assembly tree. Each leaf supplies one aligned seed fragment, and each internal vertex joins two adjacent fragment blocks. The contraction argument has the advantage of giving a matching lower bound for *every* representation, not merely constructing one efficient derivation.

3.2 There are no missing finite ranks

Theorem 3.4 (No-gap theorem for aligned rank). *If $L_n \neq \emptyset$, the ranks attained on $\text{Rec}(L)_n$ form exactly $\{1, 2, \dots, R_L(n)\}$ for some $R_L(n)$. For $n \geq 1$, $1 \leq R_L(n) \leq n$; for the active zero slice, $R_L(0) = 1$.*

Proof. Let w have rank r , and choose a shortest partition with boundaries $i_0, \dots, i_r$ and corresponding seed witnesses $u_1, \dots, u_r$. For each $1 \leq k \leq r$, form

$$v_k = w[0 : i_{k-1}] u_k[i_{k-1} : n].$$

The first $k - 1$ old fragments followed by the suffix of u_k show that $\rho_L(v_k) \leq k$. Also v_k agrees with w through boundary i_k . If $k \geq 2$ and v_k had a partition into at most $k - 1$ compatible fragments, restrict that partition to its prefix ending at i_k and discard empty intervals. Append the last $r - k$ fragments of the original partition of w . This would give w at most $(k - 1) + (r - k) = r - 1$ fragments, a contradiction. Thus v_k has rank exactly k . The case $k = 1$ follows from positive rank. There are finitely many words of length n , so the largest rank exists. The earlier singleton-fragment bound proves $R_L(n) \leq n$. $\square$

This no-gap assertion concerns aligned seed fragments, not the number of runs in an insertion mask. Its proof changes the output word and retains alignment throughout; it gives no proof of the both-singleton insertion conjecture in the repository.

3.3 Sharp stabilization bounds

Corollary 3.5. *An active slice stabilizes at the exact stage*

$$h_L(n) = \lceil \log_2 R_L(n) \rceil.$$

Empty slices stabilize at stage zero. In particular, for $n \geq 1$ all length- n behavior stabilizes by $\lceil \log_2 n \rceil$. The whole language stabilizes after finitely many generations if and only if its finite ranks are uniformly bounded. If their maximum is $M \geq 1$, the least global stage is $\lceil \log_2 M \rceil$; the empty language has stage zero.

Proof. A slice equals its hull exactly when the rank threshold 2^k covers its largest rank. If it equals its next generation, Lemma 2.2 and Theorem 3.3 imply that it already equals the hull. The global statement follows by taking the supremum of all finite ranks. $\square$

Example 3.6 (Sharpness with a regular seed). Let $L = 0^* \cup 0^*10^*$, the binary words having at most one 1. Its hull is all of $\{0, 1\}^*$, and

$$\rho_L(w) = \max(1, |w|_1), \quad L^{[k]} = \{w : |w|_1 \leq 2^k\}.$$

A compatible fragment contains at most one 1, giving the lower bound. Cutting between consecutive 1's gives the reverse bound, because each resulting fragment extends by zeros to an allowed seed of the same length. Thus $R_L(n) = n$ for $n \geq 1$, and the logarithmic slice bound is sharp. Nevertheless every global inclusion $L^{[k]} \subsetneq L^{[k+1]}$ is strict.

Corollary 3.7 (Shortest possible late word). *Every word in $L^{[k+1]} \setminus L^{[k]}$ has length at least $2^k + 1$. For every $k \geq 0$, Example 3.6 attains this lower bound with 1^{2^k+1} .*

Proof. A newly appearing word has rank greater than 2^k and finite rank at most its length. The example has rank equal to the number of 1's. $\square$

3.4 The frozen-source convention is also exactly solvable

For comparison, let $Y_0 = L$ and $Y_{k+1} = L \otimes_c Y_k$, using both children in Definition 2.1. This convention keeps one parent in the seed at every step. It has the same full closure, but a different clock.

Theorem 3.8 (Frozen-source generation). *For every $k \geq 0$,*

$$Y_k = \{w : \rho_L(w) \leq k + 1\}.$$

Thus a word of finite rank r has frozen-source depth exactly $r - 1$. Finite stabilization is again equivalent to bounded finite rank, and its least stage is $M - 1$ when the maximum finite rank is $M \geq 1$. The empty seed stabilizes at stage zero.

Proof. The assertion at $k = 0$ is the definition of rank one. If the assertion holds at k , restrict a compatible partition of the Y_k parent to the portion inherited by a child. It uses at most $k + 1$ fragments. The inherited portion of the seed parent uses at most one more. Every child therefore has rank at most $k + 2$.

Conversely, suppose w has a compatible partition into $r \leq k + 2$ pieces, with seeds $u_1, \dots, u_r$ and boundaries $i_0, \dots, i_r$. A rank-one word persists by crossing it with itself. For $r \geq 2$, form

$$v = w[0 : i_{r-2}] u_{r-1}[i_{r-2} : n].$$

It has rank at most $r - 1 \leq k + 1$ and agrees with w through boundary i_{r-1} . By induction $v \in Y_k$. Its prefix through that boundary, crossed with the suffix of $u_r \in L$, produces w . This proves the reverse inclusion. The stabilization statements follow because the thresholds $k + 1$ exhaust precisely the finite ranks. $\square$

The distinction is quantitative: the parallel clock is logarithmic in rank, whereas the frozen-source clock is linear. Both clocks detect the same boundedness property. Statements about one convention should not silently be read as finite-stage statements about the other.

4 One-step equality is undecidable for context-free inputs

We use the classical fact that universality of a context-free grammar over a fixed binary alphabet is co-r.e.-complete. Appendix A recalls the usual reduction from Post correspondence and explains the fixed-alphabet step. Here “co-r.e.” means that the negative instances can be effectively enumerated; completeness is under computable many-one reductions.

4.1 A guard construction with universal first generation

Let $\Sigma = \{0, 1\}$, let $@$ be a fresh letter, and put $\Gamma = \{0, 1, @\}$. For any $K \subseteq \Sigma^*$ define

$$G_K = (\Gamma^* \setminus 00\Sigma^*) \cup 00K. \quad (6)$$

This language contains every word involving $@$, every word of length less than two, and every word not beginning with 00 . The only potentially missing words are the words $00x$ with $x \notin K$.

Lemma 4.1 (One-cut repair). *For every $K \subseteq \{0, 1\}^*$,*

$$T(G_K) = \Gamma^*, \quad G_K = \Gamma^* \iff K = \{0, 1\}^*.$$

If K is context-free, a grammar for G_K is effectively constructible. If K is given by an NFA, an NFA for G_K has size polynomial in the input.

Proof. The second equivalence is immediate from the list of missing words. Every word already in G_K remains in $T(G_K)$. A missing word $w = 00x$ has length $n \geq 2$. Define two parents

$$u = w[0 : 1] @ w[2 : n], \quad v = @ w[1 : n].$$

Both contain the guard and hence belong to G_K , both have length n , and crossing after their first letter gives $u[0 : 1]v[1 : n] = w$. Thus every missing word is repaired in one step. The language $\Gamma^* \setminus 00\Sigma^*$ is a fixed regular language; union with the prefix extension $00K$ gives the asserted effective grammar and automaton constructions. $\square$

Theorem 4.2 (Resolution of the one-step decision question). *Given a context-free grammar over the fixed alphabet $\{0, 1, @\}$, deciding whether its language satisfies $L \otimes_c L = L$ is co-r.e.-complete. This remains true on the effective family in which the first crossover generation is the entire free monoid.*

Proof. By Lemma 4.1,

$$T(G_K) = G_K \iff G_K = \Gamma^* \iff K = \{0, 1\}^*.$$

This proves co-r.e.-hardness and the promised restriction. For membership in co-r.e., enumerate pairs of equal-length words and their cut positions. Grammar membership is decidable. Whenever both parents belong to L but their child does not, the enumeration finds a finite witness to failure of closure. Such a witness exists exactly when $T(L) \neq L$, since $L \subseteq T(L)$ always holds. $\square$

The use of 00 is not cosmetic. It ensures that every omitted word has at least two positions, allowing two distinct guard replacements and a genuinely internal repair cut. Without a length guard, a missing empty word or single-letter word cannot be created by aligned crossover from other words of the same length.

Corollary 4.3. *There is no computable function of grammar-description size bounding the length of the shortest counterexample to one-step closure for all nonclosed context-free grammars over this fixed alphabet.*

Proof. Such a bound would allow a finite exhaustive search. For each word length below the bound, enumerate all equal-length parents and cuts and use grammar membership to test the resulting violations. If no violation were found, the bound would certify closure, contradicting Theorem 4.2. $\square$

5 Finite stabilization and an exact block amplification

One-step undecidability does not by itself imply undecidability of eventual stabilization: the guarded examples above *always* stabilize by stage one. We now amplify a single omitted block into arbitrarily high rank.

Introduce a fresh separator # and put $\Delta = \{0, 1, @, \#\}$. Define

$$B_K = G_K(\#G_K)^*. \quad (7)$$

Thus a word belongs to B_K exactly when every block between successive separators belongs to G_K , including the initial and final blocks. Empty blocks are allowed because $\varepsilon \in G_K$. This construction preserves context-freeness effectively.

Theorem 5.1 (Exact rank of repeated forbidden blocks). *Suppose $x \notin K$, put $w = 00x$, and let $m = |w| \geq 2$. For $r \geq 1$ define*

$$W_r = w(\#w)^{r-1}.$$

Then

$$|W_r| = (m+1)r - 1, \quad \rho_{B_K}(W_r) = r + 1, \quad d_{B_K}(W_r) = \lceil \log_2(r+1) \rceil. \quad (8)$$

Proof. Number the r copies of w . The j th copy occupies the interval between boundaries

$$s_j = (j-1)(m+1), \quad t_j = s_j + m.$$

Consider any compatible-fragment partition of W_r . We claim it must have an internal partition boundary in the integer interval $[s_j, t_j]$ for each j .

If there is no such boundary, one partition fragment contains the entire j th copy of w . When a neighboring separator exists, the fragment also contains that separator, because a cut immediately before or after the copy would be a boundary in $[s_j, t_j]$. At the beginning or end of the whole word, the corresponding end of the seed is fixed by alignment and equal total length. Therefore every complete seed matching this fragment contains w as a whole separator-delimited block. This is impossible in B_K , because $w = 00x \notin G_K$.

The r integer intervals $[s_j, t_j]$ are pairwise disjoint: $s_{j+1} = t_j + 1$. Thus at least r distinct internal partition boundaries are needed, and there are at least $r+1$ fragments.

For the upper bound, place one cut strictly inside each copy of w , for example immediately after its first letter. These cuts give $r+1$ fragments. No resulting fragment contains a whole copy of w . For any one fragment, retain the separator positions of W_r , retain the letters of that fragment, and replace all other nonseparator positions by @. In every block at least one position has been replaced, so every block contains a guard and belongs to G_K . The resulting full word lies in B_K and matches the chosen fragment at the required position. Hence the partition is compatible and has $r+1$ pieces. The depth formula follows from Theorem 3.3. $\square$

The lower bound allows seed words with arbitrary separator placements outside the chosen fragment. It does not quietly restrict the minimization to seeds with the same global block layout. The separators *inside* an overlong fragment, together with the actual word endpoints, force the bad block.

Theorem 5.2 (Finite-stabilization undecidability). *It is undecidable, given a context-free grammar over the fixed four-letter alphabet Δ , whether its parallel self-crossover iterates eventually stabilize. This holds even under the promise that the language is either Δ^* or has strict inclusions at every generation. For each fixed $k \geq 0$, deciding $L^{[k]} = L^{[k+1]}$ is co-r.e.-complete.*

Proof. If $K = \{0, 1\}^*$, then $G_K = \Gamma^*$ and $B_K = \Delta^*$, so stabilization occurs at stage zero. Otherwise fix $x \notin K$ and apply Theorem 5.1. For every $k \geq 0$ the word W_{2^k} has rank $2^k + 1$, and therefore

$$W_{2^k} \in B_K^{[k+1]} \setminus B_K^{[k]}, \quad |W_{2^k}| = (m+1)2^k - 1. \quad (9)$$

Thus B_K satisfies the promise and stabilizes if and only if K is universal. This proves undecidability and co-r.e.-hardness for each fixed adjacent equality problem.

For a fixed k , membership in either finite generation is decidable: for a candidate word of length n , one may enumerate the finite set of seed words of length n using grammar membership, then compute crossover $k+1$ times. Enumerating words consequently semidecides a failure of the adjacent equality. Section 6 gives a much faster membership algorithm. $\square$

Corollary 5.3 (The same decision boundary with a frozen source). *Finite stabilization is also undecidable for $Y_{k+1} = L \otimes_c Y_k$, and every fixed adjacent equality problem $Y_k = Y_{k+1}$ is co-r.e.-complete for CFGs. On the nonuniversal side of the block construction, W_{k+1} has frozen-source depth $k+1$ and therefore witnesses strictness at stage k .*

Proof. Theorem 3.8 turns the exact block rank $r+1$ into depth r . Use $r = k+1$ for a fixed-stage witness. The effective reduction and the semidecision of failures are the same as in Theorem 5.2. $\square$

Proposition 5.4 (An arithmetical upper bound). *The set of context-free grammars whose self-crossover sequence eventually stabilizes belongs to Σ_2^0 and is Π_1^0 -hard. No completeness claim at level Σ_2^0 is made here.*

Proof. The condition has the form

$$\exists k \forall w \quad (w \in L^{[k+1]} \implies w \in L^{[k]}),$$

with a decidable matrix, uniformly in the grammar, k , and w . Theorem 5.2 reduces the co-r.e.-complete universality problem to this condition, proving Π_1^0 -hardness. $\square$

6 Polynomial-time rank computation for a fixed target

Global undecidability is compatible with efficient local membership. The aligned-fragment invariant makes this particularly transparent.

Proposition 6.1 (Compatibility by constrained parsing). *Let G be a context-free grammar in Chomsky normal form, of size g , and let w have length $n \geq 1$. All interval-compatibility answers for w can be computed in $O(gn^5)$ time. Consequently $\rho_{L(G)}(w)$ and its first generation, when finite, can be computed in that time. The bound is in the size after the usual polynomial-time grammar normalization.*

Proof. For $0 \leq i < j \leq n$, compatibility is equivalent to

$$L(G) \cap \Sigma^i \{w[i:j]\} \Sigma^{n-j} \neq \emptyset. \quad (10)$$

Run CYK on a length- n string whose positions inside $[i, j)$ are fixed to the corresponding letters of w and whose other positions are wildcards. A terminal production is allowed at a wildcard if it produces any alphabet letter; at a fixed position it must produce that letter. The binary CYK recursion is unchanged. Its acceptance means that some full word meeting these position constraints belongs to the language. Each query costs $O(gn^3)$, and there are $O(n^2)$ intervals.

After the queries, make a directed acyclic graph on boundaries $0, 1, \dots, n$, with an edge $i \rightarrow j$ exactly for the compatible intervals. The least number of edges from 0 to n is exactly the fragment rank. Equivalently,

$$D[0] = 0, \quad D[j] = 1 + \min\{D[i] : 0 \leq i < j, [i, j) \text{ compatible}\},$$

with the minimum of the empty set equal to ∞ . This last step takes $O(n^2)$ time. The depth is obtained from (5). $\square$

This algorithm answers membership in $L^{[k]}$ even when k is given in binary and is very large: compare the computed depth with k , rather than constructing the integer 2^k . A word has finite depth at most $\lceil \log_2 n \rceil$. The empty word is handled by a single grammar membership test. Back pointers in the dynamic program and parsing tables recover a shortest partition and its seed witnesses; balanced pairing then recovers an explicit crossover derivation.

The conservative $O(gn^5)$ bound does not assert optimality. Sharing work between wildcard parses should improve it. More importantly, no step requires a context-free description of an intermediate generation. Such a description has not been assumed to exist.

7 Regular inputs: finite-state tests and limitedness

7.1 One-step closure for deterministic automata

Let $\mathcal{A} = (Q, \Sigma, \delta, q_0, F)$ be a complete DFA with s states. The completion convention includes a rejecting sink when necessary.

Theorem 7.1 (A cubic-state violation graph). *One can decide $T(L(\mathcal{A})) = L(\mathcal{A})$ in $O(s^3|\Sigma|^2)$ time by graph reachability. If closure fails, there is a counterexample word of length at most $2s^3 - 1$.*

Proof. Use states $(p, q, r, e) \in Q^3 \times \{0, 1\}$. The first two components read the two parent words; the third reads the child; e records whether the cut has occurred. Initially all three states are q_0 and $e = 0$. There is a single possible phase change $(p, q, r, 0) \rightarrow (p, q, r, 1)$, consuming no letter.

On a child letter a , choose an arbitrary letter b for the parent whose letter is not inherited. In phase zero update

$$(p, q, r) \mapsto (\delta(p, a), \delta(q, b), \delta(r, a));$$

in phase one update

$$(p, q, r) \mapsto (\delta(p, b), \delta(q, a), \delta(r, a)).$$

A state is accepting for the violation graph when $p, q \in F$ and $r \notin F$. Every accepting path supplies two accepted equal-length parents and a rejected child, with one aligned cut. Conversely every closure violation supplies such a path, using phase changes at the endpoints when needed.

There are $2s^3$ states and $O(s^3|\Sigma|^2)$ transitions. Reachability therefore has the stated complexity. A shortest edge path to an accepting state is simple, has at most $2s^3 - 1$ edges, and consumes no more letters than that. Thus a counterexample of the asserted length exists. $\square$

Ordinary graph search finds a bounded witness; using zero–one breadth-first search, assigning zero to the phase change and one to letter transitions, finds a shortest child word. The graph tracks the two parents independently outside the inherited portions, which is essential for exact-length matching.

Theorem 7.2 (NFA complexity). *For a language given by an NFA, one-step crossover closure is PSPACE-complete, already over a fixed three-letter alphabet.*

Proof. For hardness, apply the polynomial NFA construction in Lemma 4.1 to binary NFA universality. The latter problem is PSPACE-complete; fixed-alphabet hardness is classical, and is also available under stronger automaton restrictions [6].

For membership, first eliminate epsilon transitions. Adapt the violation graph: its parent components choose individual NFA runs, while its child component is the subset of all states reachable on the child prefix. Acceptance requires accepted parent runs and a child subset disjoint from F . A state consists of two NFA states, one subset, and one phase bit, and can be stored in polynomial space. Reachability in the exponentially large implicit graph is decidable in polynomial space. PSPACE is closed under complementation, so the closure rather than violation problem has the same upper bound. $\square$

7.2 Encoding all fragment ranks by a distance automaton

A *distance automaton* is a finite nondeterministic automaton with transition costs 0 or 1. The value of a word is the minimum total cost of an accepting path, or ∞ if there is no accepting path. The classical *limitedness problem* asks whether all finite values are bounded by one integer. It is decidable in PSPACE in the explicit automaton size [5]. We use this theorem, rather than reprove it.

For a complete DFA $\mathcal{A}$ define, for $S \subseteq Q$,

$$\text{Post}(S) = \{\delta(q, a) : q \in S, a \in \Sigma\}, \quad \text{Pre}(S) = \{q : \exists a \in \Sigma, \delta(q, a) \in S\}.$$

For a target length n , set

$$U_i = \delta(q_0, \Sigma^i), \quad V_i = \{q : \delta(q, v) \in F \text{ for some } v \in \Sigma^{n-i}\}.$$

Then $U_0 = \{q_0\}$, $U_{i+1} = \text{Post}(U_i)$, $V_n = F$, and $V_i = \text{Pre}(V_{i+1})$. An interval $w[i : j]$ is compatible exactly when reading it from some state in U_i ends in V_j .

Theorem 7.3 (A finite-state representation of rank). *From an s -state complete DFA for L one can construct a distance automaton with at most $s4^s$ states, allowing a set of initial states, whose finite-word value is*

$$\text{val}(w) = \rho_L(w) - 1$$

when $\rho_L(w) < \infty$, and ∞ otherwise. In particular $\text{Rec}(L)$ is effectively regular.

Proof. A state is a triple (U, V, q) , where $U, V \subseteq Q$ and $q \in Q$. The initial states have $U = \{q_0\}$, $q = q_0$, and arbitrary V . Final states satisfy $V = F$ and $q \in F$.

When consuming the next child letter a , choose a new subset V' subject to $\text{Pre}(V') = V$ and set $U' = \text{Post}(U)$. There are two sorts of transitions:

$$\begin{aligned} (U, V, q) &\longrightarrow (U', V', \delta(q, a)) && \text{at cost 0;} \\ (U, V, q) &\longrightarrow (U', V', \delta(p, a)) && \text{at cost 1, } q \in V, p \in U. \end{aligned}$$

The second transition closes the current fragment and starts a new one just before consuming a .

For any accepting path of length n , the boundary conditions $V_n = F$ and $V_i = \text{Pre}(V_{i+1})$ force the guessed subsets to be exactly the sets V_i defined above. The U_i are forced by their forward recurrence. A completed

fragment starts in a state reachable by an arbitrary prefix of the required length and ends in a state having an arbitrary accepting suffix of the required length. It therefore extends to a genuine seed in L_n . The last fragment ends in F and has the same property.

An accepting path of cost c thus gives a compatible partition into at most $c + 1$ nonempty fragments. A reset before the first letter is redundant: $U_0 = \{q_0\}$, so deleting that reset preserves its state transition and reduces its cost. There are no resets after the last letter, because resets are combined with a consumed letter. A minimum-cost path consequently corresponds to exactly its positive-length fragments.

Conversely, given a compatible partition into r fragments, choose the states and runs of their seed witnesses. Continue at cost zero within a fragment and reset at cost one at each of its $r - 1$ internal boundaries. The accepting path has cost $r - 1$. This proves the equality of values. For ε , a cost-zero accepting initial state exists exactly when $q_0 \in F$, agreeing with the rank convention.

There are $2^s \cdot 2^s \cdot s = s4^s$ states. Forgetting the weights yields an NFA recognizing the finite-rank domain, which is $\text{Rec}(L)$ by Theorem 3.3. $\square$

Corollary 7.4 (Decidable regular-input stabilization). *Finite stabilization of parallel self-crossover, and of frozen-source crossover, is decidable for regular languages. For complete DFA input, the preceding construction and the classical limitedness algorithm give an EXPSPACE upper bound in the DFA size. This is an upper bound, not a claimed optimal complexity.*

Proof. By Corollary 3.5 and Theorem 3.8, either form of stabilization is equivalent to bounded finite rank. Theorem 7.3 converts this into limitedness of a distance automaton. Its explicit state and transition description has size $2^{\text{poly}(s+|\Sigma|)}$; polynomial space in that size is exponential space in the original description size. $\square$

[Added 30 September 2026, batch 66: Part II improves this upper bound to PSPACE, for NFA and for DFA input (Theorem 19.2), by a periodic-interior criterion instead of generic limitedness, and proves the problem PSPACE-complete for NFA input over a binary alphabet (Corollary 20.3); every stabilizing s -state NFA language has rank at most $(2s + 5)^2$ (Corollary 19.3). The exact complexity for DFA input stays open. The EXPSPACE bound stated here remains true; Part II is unrefereed and not formalized.] [Added 1 October 2026, batch 73: Part VI settles the exact complexity for DFA input: finite stabilization of an explicit complete DFA is coNP-complete, already over a binary alphabet (Theorem 67.1).]

When limitedness holds, the least stabilization stage is also computable. Indeed, each finite iterate of a regular language is effectively regular by a product-and-cut construction, and equality of regular languages is decidable. Generate iterates and test adjacent equality until the first one occurs. The limitedness decision guarantees termination of this second procedure on the bounded side. This gives an effective answer, not a practical implementation of an optimal limitedness solver.

8 Rational counting for the full closure of every CFL

The regularity conclusion of Theorem 7.3 does not extend to context-free inputs. A weaker conclusion does, and it is stronger than rationality of the ordinary length series: all commuting letter weights may be retained.

8.1 Effective semilinear position supports

Lemma 8.1. *Let $L \subseteq \Sigma^*$ be context-free. For each $a \in \Sigma$, the set*

$$E_a = \{(i, j) \in \mathbb{N}^2 : \exists u \in \Sigma^i, v \in \Sigma^j, uav \in L\} \quad (11)$$

is effectively semilinear, and hence effectively definable in Presburger arithmetic.

Proof. Take two tagged copies of the alphabet and a new marker $\bullet$ that is mapped to a . The inverse image of L under the letter-forgetting homomorphism is context-free. Intersect it with the regular language “left-tagged letters, then $\bullet$, then right-tagged letters.” Its Parikh image is effectively semilinear by the effective form of Parikh’s theorem [3]. Sum the counts of left-tagged letters, sum the counts of right-tagged letters, and discard the marker coordinate. The resulting linear image is exactly E_a . $\square$

For $0 \leq i < n$, the letters allowed at position $i + 1$ of a hull word are

$$P_{n,i+1}(L) = \{a \in \Sigma : (i, n - 1 - i) \in E_a\}.$$

Therefore, for each $S \subseteq \Sigma$, the predicate $P_{n,i+1}(L) = S$ is a Boolean combination of effective Presburger predicates. Define the number of positions of each type by

$$A_S(n) = \#\{i : 0 \leq i < n, P_{n,i+1}(L) = S\}. \quad (12)$$

Lemma 8.2 (Eventually affine counts on residue classes). *There are an effectively obtainable period $p \geq 1$ and, for each residue class, a starting length N_r in that class such that simultaneously for all $S \subseteq \Sigma$,*

$$A_S(N_r + pt) = b_{S,r} + d_{S,r}t \quad (t \geq 0), \quad (13)$$

with $b_{S,r}, d_{S,r} \in \mathbb{N}$. Each such progression can also be chosen either entirely active ($L_n \neq \emptyset$) or entirely inactive.

Proof. Each A_S is a Presburger counting function. Woods’s counting theorem implies that it is a piecewise quasi-polynomial [4, Theorem 1.10]. In one parameter, beyond a finite threshold the polyhedral partition has only unbounded rays, so the function is an eventual quasi-polynomial. Since $0 \leq A_S(n) \leq n$, each constituent polynomial has degree at most one. Take a common period and a sufficiently large starting length in each residue class for all the finitely many subsets S .

The value at the starting length is an integer $b_{S,r} \geq 0$, and the constant first difference along the progression is an integer $d_{S,r}$. A negative first difference would eventually make the count negative, so $d_{S,r} \geq 0$. These procedures are effective: Presburger quantifier elimination and the counting construction obtain a representation, and all subsequent finite refinements are effective. Appendix B also explains the one-parameter affine conclusion directly.

The length set of a context-free language is effectively semilinear by Parikh’s theorem. A semilinear subset of $\mathbb{N}$ is eventually periodic. Refine the same period and thresholds so that activity is constant on each progression. This deals explicitly with empty slices. $\square$

8.2 A rational commutative generating function

For commuting variables $\mathbf{x} = (x_a)_{a \in \Sigma}$ put

$$\text{wt}(w) = \prod_{a \in \Sigma} x_a^{|w|_a}, \quad \Phi_L(z, \mathbf{x}) = \sum_{w \in \text{Rec}(L)} z^{|w|} \text{wt}(w).$$

The series counts distinct words, not derivations or seed representations.

Theorem 8.3 (Rational commutative growth). *For every context-free language L , the series $\Phi_L(z, \mathbf{x})$ is an effectively computable rational function. More precisely, it has a representation*

$$\Phi_L(z, \mathbf{x}) = P(z, \mathbf{x}) + \sum_{r \in \mathcal{R}} \frac{z^{N_r} C_r(\mathbf{x})}{1 - z^p D_r(\mathbf{x})}, \quad (14)$$

where P is a polynomial, $\mathcal{R}$ indexes the eventually active residue classes, and C_r, D_r are products of powers of the nonzero linear forms $\sum_{a \in S} x_a$ for $\emptyset \neq S \subseteq \Sigma$. In particular, along each sufficiently late active arithmetic progression,

$$|\text{Rec}(L)_{N_r + pt}| = C_r(\mathbf{1}) D_r(\mathbf{1})^t, \quad (15)$$

where both constants are positive integers. Inactive progressions have count zero.

Proof. On an active slice the choices of letters at different positions are independent. Hence

$$\begin{aligned} \sum_{w \in \text{Rec}(L)_n} \text{wt}(w) &= \prod_{i=1}^n \left(\sum_{a \in P_{n,i}(L)} x_a \right) \\ &= \prod_{\emptyset \neq S \subseteq \Sigma} \left(\sum_{a \in S} x_a \right)^{A_S(n)}. \end{aligned} \quad (16)$$

Insert (13). For $n = N_r + pt$ this becomes $C_r(\mathbf{x}) D_r(\mathbf{x})^t$, with

$$C_r(\mathbf{x}) = \prod_{S \neq \emptyset} \left(\sum_{a \in S} x_a \right)^{b_{S,r}}, \quad D_r(\mathbf{x}) = \prod_{S \neq \emptyset} \left(\sum_{a \in S} x_a \right)^{d_{S,r}}.$$

Summing this geometric series in t gives its summand in (14). The finitely many lengths before the chosen starts, including length zero, supply P . Their coordinate sets and activity are decidable, so this polynomial is effective. Specializing each x_a to one yields (15). $\square$

Corollary 8.4. *The ordinary length generating function of $\text{Rec}(L)$ is effectively rational. Its coefficients satisfy an effectively computable constant-coefficient linear recurrence. If $\text{Rec}(L)$ is infinite, its upper exponential growth rate is*

$$\limsup_{n \rightarrow \infty} |\text{Rec}(L)_n|^{1/n} = \max_{r \in \mathcal{R}} D_r(\mathbf{1})^{1/p}.$$

Proof. Rationality and the recurrence follow by setting $\mathbf{x} = \mathbf{1}$. On an active progression the n th root of $C_r(\mathbf{1}) D_r(\mathbf{1})^t$, with $n = N_r + pt$, tends to $D_r(\mathbf{1})^{1/p}$. There are finitely many progressions. Infinitude ensures that at least one is active. $\square$

This is not a rationality theorem for the noncommutative word series and not a regularity theorem for the hull. The generating function has forgotten *where* each coordinate type occurs. Its rational form reflects the eventually affine *number* of positions of each type. The next family shows why that distinction is substantive.

9 A linear context-free seed with a non-context-free closure

9.1 Definition, grammar, and position geometry

Fix an integer $q \geq 3$ and define a binary seed

$$\begin{aligned} S_q = 0^* \\ \cup \{0^i 1 0^j : j \geq (q-1)i + q - 1\} \\ \cup \{0^i 1 0^j : i \geq (q-1)j + q - 1\}. \end{aligned} \quad (17)$$

Every seed has at most one 1. The asymmetric inequalities allow that 1 only in an initial or final portion of the word.

Lemma 9.1. *The language S_q is linear context-free. If $n = qm + r$ with $0 \leq r < q$, then the coordinate hull permits 1 exactly in the first m and last m positions.*

Proof. For the second language in (17), use

$$A \rightarrow 0A0^{q-1} \mid 10^{q-1}Z, \quad Z \rightarrow 0Z \mid \varepsilon.$$

After i recursive uses, the trailing zero count is $(q-1)i + q - 1 + t$, $t \geq 0$, as required. Reversing the productions gives a linear grammar for the third language. Their union with 0^* is linear: each production still has at most one nonterminal.

For a word of length $n = i + 1 + j$, the first inequality is equivalent to $n \geq q(i + 1)$. Thus the allowed position $i + 1$ is at most $\lfloor n/q \rfloor = m$. The second inequality gives the symmetric last m positions. The all-zero seed allows 0 everywhere. $\square$

Write $K_q = \text{Rec}(S_q)$. Lemma 9.1 gives the exact description

$$K_q = \bigcup_{\substack{m \geq 0 \\ 0 \leq r < q}} \{u 0^{(q-2)m+r} v : u, v \in \{0, 1\}^m\}. \quad (18)$$

For example, at length 10 and $q = 3$ the first and last three positions are free, while the middle four positions are zero. There are $2^6 = 64$ hull words, but only seven seed words.

Theorem 9.2 (The closure is not context-free). *For every $q \geq 3$, the language K_q is not context-free, although its seed S_q is linear context-free.*

Proof. Intersect K_q with the regular language $1^+ 0^+ 1^+$. A word $1^a 0^b 1^c$ with $a, b, c \geq 1$ lies in the intersection exactly when

$$(q-1)a \leq b + c, \quad (q-1)c \leq a + b. \quad (19)$$

Indeed, the left block of ones fits in the first $\lfloor n/q \rfloor$ positions exactly when $qa \leq n = a + b + c$, and similarly for the right block.

Suppose this intersection were context-free, with pumping length p . The word

$$z = 1^p 0^{(q-2)p} 1^p$$

satisfies both inequalities at equality. For any admissible pumping factorization, the two pumped portions are contained in one window of length at most p , and their combined length is positive. This window cannot meet both outer blocks of ones, because the intervening zero block has length at least p .

Assume, as the pumping property would require, that pumping both down to zero copies and up to two copies stays in the intersection. If either pumped word fails to have three nonempty blocks in the required order, this already is a contradiction. Otherwise let $\delta_a, \delta_b, \delta_c$ be the numbers of letters removed from the three original blocks by pumping down. The up-pumped word adds these same counts. Apply (19) to both pumped words and use equality for z . Opposite signs force

$$(q-1)\delta_a = \delta_b + \delta_c, \quad (q-1)\delta_c = \delta_a + \delta_b.$$

Subtracting gives $\delta_a = \delta_c$, and then $\delta_b = (q-2)\delta_a$. The window misses at least one outer block, so one of δ_a, δ_c is zero. All three counts are therefore zero, contradicting the positive combined length of the pumped portions. The intersection is not context-free, hence neither is K_q . $\square$

9.2 Every rank, every generation, and exact counts

Theorem 9.3. *For $w \in K_q$,*

$$\rho_{S_q}(w) = \max(1, |w|_1). \quad (20)$$

Consequently

$$S_q^{[k]} = K_q \cap \{w : |w|_1 \leq 2^k\}. \quad (21)$$

For $n = qm + r$, $0 \leq r < q$,

$$|(K_q)_n| = 4^m, \quad (22)$$

$$|(S_q^{[k]})_n| = \sum_{j=0}^{2^k} \binom{2m}{j}, \quad (23)$$

where $\binom{2m}{j} = 0$ for $j > 2m$. The exact slice stabilization stage is $\lceil \log_2 \max(1, 2m) \rceil$.

Proof. Any fragment matching a seed contains at most one 1, so rank is at least $\max(1, |w|_1)$. Partition w between its consecutive ones. Each fragment then has at most one 1; if it has one, that position is seed-supported by Lemma 9.1. Fill all other positions of the full word with zero to extend the fragment to a seed. A fragment with no ones extends to the all-zero seed. This proves (20), including the all-zero and empty cases under the stated convention.

There are $2m$ free binary coordinates. Choosing which j of them are one gives $\binom{2m}{j}$ words. Sum over all j for the hull and over $j \leq 2^k$ for generation k . The maximum rank is $\max(1, 2m)$, giving the last assertion. $\square$

Marking each 1 by a variable x yields

$$\sum_{w \in K_q} z^{|w|} x^{|w|_1} = \frac{1 + z + \cdots + z^{q-1}}{1 - (1+x)^2 z^q}. \quad (24)$$

In particular,

$$\sum_{n \geq 0} |(K_q)_n| z^n = \frac{1 + z + \cdots + z^{q-1}}{1 - 4z^q}. \quad (25)$$

The rational series is exceptionally simple, even though the underlying language is not context-free.

9.3 Sharp recurrence orders of the finite generations

Theorem 9.4 (Exponential growth of minimal recurrence order). *Fix $q \geq 3$ and $k \geq 0$, and write $D = 2^k$. The ordinary length generating function of $S_q^{[k]}$, in reduced form over $\mathbb{Q}$, has denominator*

$$(1 - z)(1 - z^q)^D. \quad (26)$$

Its minimal eventual constant-coefficient recurrence has order $qD + 1 = q2^k + 1$. For each fixed k , its coefficients grow polynomially of degree D , while the full hull has exponential growth rate $4^{1/q}$.

Proof. The polynomial

$$b_D(m) = \sum_{j=0}^D \binom{2m}{j}$$

has degree exactly D and leading coefficient $2^D/D!$. Interpreting each binomial as a polynomial in m agrees with the combinatorial convention at all $m \in \mathbb{N}$, since the falling factorial vanishes when $j > 2m$. Consequently

$$\sum_{m \geq 0} b_D(m)t^m = \frac{Q_D(t)}{(1-t)^{D+1}}, \quad \deg Q_D \leq D, \quad Q_D(1) = 2^D \neq 0.$$

One proof is to expand $b_D(m)$ in the binomial-polynomial basis; the coefficient of $\binom{m}{D}$ is $D! \cdot 2^D/D! = 2^D$, and only that term contributes to the numerator at $t = 1$.

Lengths $qm, \dots, qm + q - 1$ have the same coefficient. Thus

$$\begin{aligned} \sum_{n \geq 0} |(S_q^{[k]})_n| z^n &= (1 + z + \dots + z^{q-1}) \frac{Q_D(z^q)}{(1 - z^q)^{D+1}} \\ &= \frac{Q_D(z^q)}{(1 - z)(1 - z^q)^D}. \end{aligned}$$

All roots of the displayed denominator satisfy $z^q = 1$, where its numerator is $Q_D(1) \neq 0$. Hence it is reduced and its degree $qD + 1$ is the minimal recurrence order. Finally $b_D(\lfloor n/q \rfloor) \sim (2^D/(D!q^D))n^D$, whereas $|(K_q)_n| = 4^{\lfloor n/q \rfloor}$ has exponential rate $4^{1/q}$. $\square$

For the frozen-source sequence of Theorem 3.8, the same calculation holds with $D = k + 1$ rather than $D = 2^k$. Its generation- k minimal recurrence order is therefore $q(k + 1) + 1$.

The finite-generation and full-closure limits therefore do not commute at the level of exponential growth. Every fixed generation has growth rate one, but their increasing union has larger growth rate. This supplies a concrete warning against extrapolating an infinite closure from a few rational finite-stage enumerators.

q	generation k	rank threshold 2^k	polynomial growth degree	recurrence order
3	0	1	1	4
3	1	2	2	7
3	2	4	4	13
3	3	8	8	25
4	2	4	4	17
5	3	8	8	41

10 Verification, reproducibility, and proof boundaries

10.1 Executed exact audits

The accompanying standard-library Python program performs independent finite checks using explicit source sets, direct crossover, interval dynamic programming, finite-state paths, and integer counting. The delivered run completed **1,766,639 checks**, all passing. The JSON receipt records the detailed counts; these are finite audits, not universal proof certificates.

Audit	Executed scope
All finite binary slices	All 65,814 seed subsets at lengths 0, 1, 2, 3, 4; 1,050,698 target-rank queries. Checks exact contraction, coordinate hull, every finite rank, generation thresholds, frozen-source generations, and stabilization.
All two-state binary DFAs	All 64 complete transition/final-state choices with fixed initial state; 8,128 target queries at lengths 0 through 6. Compares source enumeration, DFA fragment rank, and the distance-automaton cost.
Guard repair	All 255 omitted binary patterns 00x of lengths 2 through 9; checks the actual same-length parent repair.
Block amplification	Ranks and distances of repeated forbidden block 00 for 1 through 64 blocks. Independent enumeration over the full four-letter alphabet at lengths 2, 5, 8.
The linear-seed family	$q = 3, \dots, 7$, all binary target words through length 12; exact ranks, weighted histograms, and finite-stage counts.
Recurrence checks	25 parameter pairs, $q = 3, \dots, 7$ and $k = 0, \dots, 4$; integer convolution with the proved denominator and nonzero top finite differences.

For the distance-automaton audit, the accepting boundary condition forces a unique sequence of right subsets V_i . The program computes that sequence backward rather than enumerate unsuccessful guesses. It then evaluates the cost transitions separately from interval-partition dynamic programming. This checks the accepted-path cost semantics; it does not implement or test the general limitedness algorithm imported from the literature.

The block test uses the regular specialization $G = \{0, 1, @\}^* \setminus \{00\}$. This does not prove the grammar undecidability reduction. It checks the combinatorial amplification while allowing arbitrary source separator placements, including the placements that could invalidate a careless lower-bound proof.

10.2 Files and rebuilding

The main source is standalone and has its bibliography internally. The package layout is

```
ProveIt_Constrained_Crossover/
  article.tex
  article.pdf
  README.md
  Makefile
  code/verify.py
  data/verification.json
  data/verification.txt
  notes/provenance.md
  notes/proof-audit.md
```

From the package directory, run

```
python3 code/verify.py
latexmk -pdf -interaction=nonstopmode -halt-on-error article.tex
```

The verifier uses no external packages and no floating-point mathematics; the only floating-point output is elapsed wall-clock time. Assertions are implemented as explicit exception-raising checks and remain active under `python -O`. The LaTeX build uses commonly distributed packages, including `newtxtext` and `newtxmath`.

[write] *Shipped layout*. In the collection the files live at the report root,

```
SetTheory/Cardinals/docs/reports/
automata-and-formal-languages/constrained-crossover-closure/.
```

There, `notes/provenance.md` and `notes/proof-audit.md` are shipped as `provenance.md` and `proof-audit.md`, and `Makefile` as `code/Makefile`; the other paths are unchanged, and all these files are byte-identical to the delivery. The delivered README and 23-page PDF are not shipped: the report's README replaces the former, and `article.pdf` is a build of this text. The verifier writes `../data/verification.json` relative to its own directory, so run in place it overwrites the shipped receipt; the README gives commands for a copy. The build now also uses the `shuffle` package, for the sibling report's operation in Section 1.4. [Added 30 September 2026, batch 66: this section describes Part I's files. Part II's files carry the prefix `02-pspace-stabilization-` and Part III's the prefix `03-closure-classification-`; Sections 23 and 45 list them. Both of their verifiers also write `data/verification.json` by default, that is, over Part I's receipt, and Part III's verifier and Part II's example program import sibling modules under their delivered names, so they are rerun only on copies, as the README describes.] [Added 1 October 2026, batch 71: Part IV's files carry the prefix `04-quadratic-finite-rank-` (Section 54); its checker prints its record to standard output and writes no file.]

10.3 What was not verified or claimed

No Lean or Rocq kernel has checked the proofs. The external Parikh, Presburger, and distance-limitedness theorems are explicitly cited mathematical inputs, not imported formal declarations. No general CFG-to-Presburger compiler, limitedness solver, or full symbolic rational-enumerator implementation is included. The polynomial-time grammar rank algorithm is proved but the executable reference verifier specializes to finite seeds and DFAs.

The novelty search inspected the cited Hughes version, the relevant ProveIt source and collection catalogue, and targeted literature on crossover, universality, effective Parikh images, Presburger counting, and distance automata. It did not identify a prior resolution of the specific constrained one-step question or an equivalent package of the quantitative results. Search coverage is not an exhaustive bibliography and is not a proof of priority. In particular, the product-hull observation and possible equivalent rank formulations deserve review by specialists in recombination systems.

[write] Placement in ProveIt, a repository with Lean and Rocq developments, confers no formal status on this report, and none of its statements is formalized. ProveIt has no formal development of crossover, aligned rank, context-free grammars, Post correspondence, Parikh images, Presburger counting or distance-automaton limitedness. The nearest formal neighbour is `Logic/PresburgerArithmetic`, whose Lean development proves Cooper quantifier elimination for Presburger arithmetic over $\mathbb{Z}$ (`PresburgerArithmetic.Formula.holds_iff_quantifierEliminate`

and the decision procedure `PresburgerArithmetic.Formula.presburgerArithmetic_decidable`, in `Lean/PresburgerArithmetic/Decision.lean`), with an independent Rocq proof of the one-variable Cooper step. The effectivity argument of Lemma 8.2 and Appendix B invokes quantifier elimination of this kind but is not connected to that development, and the counting theorem it needs is not formalized anywhere in ProveIt. The Lean plan of Section 11.1 is a plan, not an executed development.

11 Further research questions and a formalization plan

The following are proposed research tasks relative to the present results. Only the Hughes question identified in Section 1 is assigned a specific published open-problem provenance. The rest are not asserted to be newly discovered or globally open without further literature review.

Question 11.1 (Exact arithmetical complexity). Is finite stabilization of a context-free seed Σ_2^0 -complete? The present proof gives a Σ_2^0 upper bound and Π_1^0 -hardness, not a matching lower bound. A stronger construction must encode a genuinely existential bound rather than just universality of one grammar.

[Added 30 September 2026, batch 66: Still open. Part II restates it (Question 24.6) and explains why its regular-input analysis does not reach it.]

Question 11.2 (Sharp regular-input complexity). What is the optimal complexity of deciding finite stabilization for DFA and NFA inputs? Can the $s4^s$ distance-automaton construction be represented more economically, or can its special two-sided subset structure yield a better algorithm than generic limitedness? Matching hardness results would determine whether the EXPSPACE upper bound is loose.

[Added 30 September 2026, batch 66: Answered for NFAs in Part II: finite stabilization of NFA input is PSPACE-complete, already over a binary alphabet and for factorial languages with universal hull (Corollary 20.3); the special two-sided structure does give a better algorithm than generic limitedness (Theorem 18.3). For DFAs, Part II gives a PSPACE upper bound (Theorem 19.2), so the EXPSPACE bound is loose, and a safe-core criterion (Theorem 21.2); the exact DFA complexity stays open (Question 24.1). Part III's Theorem 39.2 gives coNP-completeness of a depth threshold for succinct ray inputs, a different input format.] [Added 1 October 2026, batch 73: Part VI settles the exact complexity for DFA input: finite stabilization of an explicit complete DFA is coNP-complete, already over a binary alphabet (Theorem 67.1); so for DFAs the PSPACE bound is not tight unless coNP = PSPACE. The question is now answered for both input types.]

Question 11.3 (Two-letter undecidability). Can Theorems 4.2 and 5.2 be proved over a binary alphabet? Fixed-length coding does not automatically preserve one-step crossover: a cut inside a codeword can create spurious children. A successful binary encoding must control such internal cuts, not merely preserve language universality.

Question 11.4 (Grammar classes preserving context-freeness). Which context-free seeds have a context-free coordinate hull? The linear counterexample rules out linearity as a sufficient condition. Restrictions on the position-support sets E_a , rather than just grammar ambiguity or linearity, may provide a useful criterion. Deterministic and visibly pushdown subclasses are natural places to investigate.

[Added 30 September 2026, batch 66: Partly answered in Part III. For every context-free seed, regularity of the hull is decidable, by finite row type of the supports E_a , with a DFA or an infinite obstruction (Theorems 33.1 and 33.2). For binary one-marker seeds $L_R = 0^* \cup \{0^i 10^j : (i, j) \in R\}$ with semilinear R and finite stabilization, the hull is regular, linear context-free and non-regular, or not context-free, according

to a test on the velocities of the arithmetic rays of R (Theorem 37.1, Corollary 37.2), a criterion of exactly the suggested kind. The general question, including nonsparse masks (Question 42.1) and the deterministic and visibly pushdown subclasses, stays open.]

Question 11.5 (Faster target-word algorithms). Can the $O(gn^5)$ wildcard-parsing method be reduced to one shared chart computation of cubic or near-cubic complexity? A useful implementation should return a shortest fragment partition and full seed witnesses, not just a membership bit, so that crossover derivations can be independently checked.

Question 11.6 (Which rational counting series occur?). Characterize the commutative series realized by context-free recombination closures. Theorem 8.3 forces eventual products of linear forms along arithmetic progressions. This is more restrictive than arbitrary rational nonnegative series. Which compatibility constraints on the semilinear position supports are necessary and sufficient for realization?

[Added 30 September 2026, batch 66: Part III adds examples, not a characterization: the series $(1 + z + 4z^2)/(1 - z^3)$ of a non-context-free closure with at most four words per length (Theorem 40.1) and the family (76).]

Question 11.7 (Restricted cuts and multipoint crossover). For cut positions restricted by a regular or Presburger condition, what replaces aligned fragment rank? For crossover with several cuts between the same two parents, different fragments are coupled by parent reuse; simply replacing 2^k by a larger exponential is not justified without a new proof.

[Added 30 September 2026, batch 66: Part II's Question 24.7 asks the same for its periodic dichotomy; neither added Part answers it.]

Question 11.8 (Quantitative transition windows). For S_q , the proportion reached by generation k at length $qm + r$ is $4^{-m} \sum_{j \leq 2^k} \binom{2m}{j}$. Study uniform asymptotics when k grows with m , rather than keeping k fixed. The explicit binomial law gives a tractable test case for crossover depth distributions in less rigid seeds.

Question 11.9 (Sharp automaton witness bounds). Determine the best possible shortest-violation bound for an s -state DFA. The $2s^3 - 1$ bound comes from a direct triple-state graph. It may not be sharp because the three components share one transition structure. Explicit lower bound families and state-complexity analyses would be informative.

[write] *Bearing on Hughes's Sections 13–14*. Hughes asks, for his insertion and shuffle iterations, for the exact arithmetical complexity of finite convergence, for restricted language families on which it becomes decidable, for whether the least stabilization index can be computed, and for a general “rank condition” behind his separator proofs; his Section 13 asks for which monotone operators eventual reaching of a fixed point is decidable. For constrained self-crossover, which is monotone and inflationary (Lemma 2.2), the results above give the following counterparts. Single-step and each fixed adjacent equality are co-r.e.-complete for grammars (Theorems 4.2 and 5.2); eventual stabilization is in Σ_2^0 and Π_1^0 -hard (Proposition 5.4), and Question 11.1 asks for the exact level. Regular inputs make stabilization decidable, and the least stage is then computable (Corollary 7.4 and the paragraph after it); every slice stabilizes by stage $\lceil \log_2 n \rceil$ (Corollary 3.5). Aligned-fragment rank with the separator amplification of Theorem 5.1 is a rank-based separator argument for this one operation, not the general condition Hughes asks for. None of this answers his questions for insertion or shuffle. [Added 30 September 2026, batch 66: For regular seeds Part II locates finite stabilization exactly (PSPACE-complete for NFA input) and bounds the least stage (Corollary 19.3); Part III decides it, and computes the least stage, for semilinear one-marker seeds (Theorems 35.1 and 38.2). These too concern crossover only.]

11.1 A realistic Lean development

A formal development can begin with fixed-length words as functions $\text{Fin}(n) \rightarrow \Sigma$, avoiding casts between lists of different lengths. A cut is a boundary $i \leq n$, and its child chooses the first parent before i and the second afterwards. A compatible partition is a strictly increasing finite boundary list, with one seed witness per interval.

The first milestones should be, in order, the two fragment-transfer lemmas used in Theorem 3.2, the exact iterate formula, the no-gap prefix-replacement argument, and the block-boundary charging lemma. These are finite combinatorial statements independent of grammars and automata, and would establish the core trusted interface.

A second layer can define an executable compatibility matrix and prove that the boundary shortest-path program computes the least rank. The finite reference tests can then be replaced by proved checkers over this interface. For the decision-theoretic reductions, a separate grammar layer must certify regular union, prefixing, concatenation, and star constructions and connect them to a formally stated universality theorem. The present paper does not claim that such a bridge already exists in ProveIt.

The counting theorem requires substantially different infrastructure: effective semilinearity of the marked grammar, Boolean Presburger operations, and the affine residue-class counting lemma. Its dependencies should be tracked separately rather than presenting the whole article as formalized because its finite combinatorics has been checked.

12 Conclusion

Constrained crossover is length-preserving, but length preservation does not make its global grammar problems decidable. A one-letter guard repairs any suitably padded omission in one generation, settling the specific one-step question. Repeated separator-delimited omissions then give rank exactly one more than the number of bad blocks, separating local logarithmic convergence from global finite stabilization.

The aligned-fragment invariant unifies the structural and algorithmic results. It computes exact depth, forces gap-free rank spectra, provides a polynomial local membership procedure, and turns regular-input stabilization into distance-automaton limitedness. At the counting level, semilinear coordinate supports imply rational commutative growth for every context-free seed, but not context-freeness of the generated language. The explicit linear family demonstrates this separation and gives exact recurrence orders at every finite generation.

The principal proofs stand independently of the finite experiments. The experiments supply a reproducible adversarial audit of their finite mechanisms. Independent proof review, priority checking, and formalization are the next steps before treating the draft as a settled research publication.

[Added 30 September 2026, batch 66: Parts II and III continue this account: the regular-input stabilization problem is PSPACE-complete for NFAs, with a quadratic rank bound replacing the generic limitedness bound, and the full closure of a context-free seed has a decidable regularity problem, while sparse one-marker seeds can lose context-freeness after a single generation. Both are equally unrefereed and unformalized.] [Added 1 October 2026, batch 71: Part IV shows that Part II's quadratic rank bound has the right order for NFAs: some s -state NFA has finite maximal rank $(s - 3)^2$. It is equally unrefereed and unformalized.] [Added 1 October 2026, batch 73: Part V shows that on each residue the maximal rank of a nonstabilizing regular seed grows with a rational, computable slope, and lowers the quadratic rank bound to $2(s - 1)^2 + 3$. It is equally unrefereed and unformalized.] [Added 1 October 2026, batch 73: Part VI shows that finite stabilization of an explicit complete DFA is coNP-complete and that the largest finite rank of an s -state complete binary DFA is

$s^2 + O(s)$, with leading constant one. It is equally refereed and unformalized.] [Added 1 October 2026, batch 73: Part VII shows that the same holds for NFAs, $F(s) = s^2 + O(s)$ over every fixed finite alphabet, and computes $F(s)$ for $s \leq 4$, with computer assistance. It is equally refereed and unformalized.]

Part II. Finite stabilization of regular seeds is PSPACE-complete

[write, batch 66] Sections 1–12 and Appendices A–D form the original report (Part I) of 29 September 2026, unchanged except for dated notes that point to the additions; its appendices are printed after Part III. Part II was added on 30 September 2026 in batch 66 of ProveIt’s incoming-report intake, from one manuscript written against Part I as printed above. For regular seeds it decides finite stabilization of constrained self-crossover in polynomial space, proves the problem PSPACE-complete for NFA input over a binary alphabet, and bounds the rank quadratically in the stabilizing case. It thereby answers the NFA half of Part I’s Question 11.2 and improves the EXPSPACE bound of Corollary 7.4 to PSPACE; the exact complexity for DFA input stays open. Section 13 records the provenance and fixes the notation of the addition. [Added 1 October 2026, batch 73: the DFA complexity is coNP-complete (Part VI).]

13 [write, batch 66] Provenance and notation of Part II

13.1 One manuscript, one addition

Part II is batch-66 manuscript 01, *Finite Crossover Stabilization Is PSPACE-Complete: periodic interior universality, quadratic rank bounds, and linear-growth obstructions*, dated 30 September 2026 (23 US-letter pages as delivered). It is an AI-assisted research draft: its title page and PDF metadata name ChatGPT as the drafting assistant and Vladimir Reshetnikov as the person it was prepared for. It arrived as the archive `ProveIt_Crossover_Stabilization.zip` (inner directory of the same name) in ProveIt commit 9ad899cbe and was placed in commit 4fee1cd07. It was written against ProveIt commit b8b0fa218 (in full b8b0fa2184a044d46ce9ed0f25f88d7bb60fa042), at which Part I’s `article.tex` had the Git blob eb4f427f; that blob is unchanged at the placement commit, so every statement the manuscript makes about “the inspected report” refers to Part I as printed above. Its code, data and audit notes are shipped beside Part I’s files with the prefix `02-pspace-stabilization-` (Part I’s unprefixes files count as 01); its manuscript, delivery README and PDF are not shipped and survive in the arrival commit. Section 23 and the report’s README list the shipped names.

Every result, proof, example, remark, figure, table, limitation and question of the manuscript is printed here, in its order. Its title page, abstract and status box became this section; its Sections 1–13 are Sections 14–26; its two appendices became Sections 27 and 28, because the report’s appendices belong to Part I. Its own reproof of Part I’s rank calculus (Proposition 15.2) is printed as a marked second proof of results stated once, in Part I; Part III reproves the coordinate hull once more, by a third route (Lemma 31.3). Text marked [write, batch 66] was added when the manuscript was written into this report; no theorem was strengthened or weakened.

13.2 The manuscript’s summary and status

Abstract, as delivered (with “the inspected ProveIt crossover report” now Part I). For a regular language, does repeated equal-length, aligned one-point self-crossover reach its full closure after finitely many generations? We prove that this problem is PSPACE-complete for NFA input, already over a binary alphabet and even on factorial languages whose coordinate hull is universal. The upper bound also applies to DFA input, improving the EXPSPACE upper bound of Part I (Corollary 7.4).

The proof replaces generic distance-automaton limitedness by a finite periodic-interior universality test.

Either every compatible interior word is one aligned seed fragment, or one forbidden word can be repeated to force linearly many fragments. In the first case, for an s -state NFA, every word in the closure has rank at most $(2s + 5)^2$, so parallel stabilization occurs by $\lceil \log_2(2s + 5)^2 \rceil$ generations. In the second case the maximum rank is linear along every sufficiently large length in at least one arithmetic progression. We also give a deterministic safe-core criterion, a polynomial-in-period obstruction bound for DFAs, and a logarithmically sharp generation lower family. Complete proofs, reproducible finite audits, and a research agenda accompany the article.

Status and scope, as delivered. Unrefereed mathematical research draft, not a Lean or Rocq development. The NFA complexity question is settled here; the exact DFA complexity and the optimal rank bound are not. Priority is not asserted: the literature audit is targeted, and the claimed extension is measured against a pinned repository revision. [Added 1 October 2026, batch 71: the order of the optimal rank bound has since been settled for NFAs in Part IV: it is quadratic (Theorem 47.1); its constant, and the order for DFAs, are not.] [Added 1 October 2026, batch 73: Part VI settles both DFA questions: the exact DFA complexity is coNP-complete (Theorem 67.1), and the order for complete DFAs is quadratic with leading constant one (Theorem 67.2).] [Added 1 October 2026, batch 73: The constant of the optimal rank bound for NFAs is one as well (Part VII, Theorem 76.2).]

13.3 Notation beside Part I and Part III

Part II uses Part I's operation $\otimes_c$, crossover map $T(L) = L \otimes_c L$, parallel iterates $L^{[k]}$, frozen-source iterates Y_k , aligned-fragment rank ρ_L and coordinate hull $\text{Rec}(L)$ (Definitions 2.1, 2.4 and 3.1). Three symbols were renamed when the manuscript was written in: its coordinate hull $\mathcal{H}(L)$ is printed as Part I's $\text{Rec}(L)$ (the same set, with the same length-zero convention), and its delimiter languages $G(K)$, $B(K)$ of Section 20 are printed as $\widehat{G}(K)$, $\widehat{B}(K)$, because they are *not* Part I's G_K , B_K . No normalization changed. The table lists the other symbols that could be misread.

In Part II	Meaning in Part II	Not to be confused with
$Y_{k+1} = Y_k \otimes_c L$	frozen-source iterates	Part I writes $L \otimes_c Y_k$; since $\otimes_c$ returns both children, the two are the same set
$\widehat{G}(K), \widehat{B}(K)$	regular delimiter languages over $\{0, 1, @ \}$ and $\{0, 1, @, \# \}$ (Proposition 20.4)	Part I's guard and block languages G_K, B_K (6), (7), which are different sets
$M_L(n)$	largest rank on the length- n hull slice; 0 on an empty slice	the same as Part I's $R_L(n)$ (Theorem 3.4) on active slices
$R(L)$	supremum of all finite ranks, with 1 adjoined	Part I's $R_L(n)$ (one slice); Part III's $R \subseteq \mathbb{N}^2$ (a marker mask)
P_i, R_j	forward and backward exact-length state sets of an NFA (32)	Part I's coordinate sets $P_{n,i}(L)$; Part I's U_i, V_i (Section 7) are the DFA case P_i, R_{n-i}
$A_{n,i}$	letters allowed at zero-based position i (35)	Part I's $P_{n,i+1}(L)$ (one-based); Part III's $A_L(n, i)$ (zero-based, the same set)
$S_{n,i}, S_{r,a}, C_{r,a}$	viable state sets and stable alphabets	Part I's seed family S_q and the factors $C_r(\mathbf{x})$ of Theorem 8.3
m, T_s, p_s, B_s, K_s	the constants (27), $K_s = \lceil \log_2 B_s \rceil$	the operator $T(L)$; Part I's B_K, K, K_q ; Part III's gap-ray languages $\text{Gap}_s(a, d)$
$\mathcal{R}$	set of live residues	Part I's index set $\mathcal{R}$ of eventually active residue classes (Theorem 8.3), a similar notion for context-free seeds
L_m	the one-1 family of Theorem 22.1	Part III's family L_k^{fam} and seeds L_R ; $L_m = L_R \setminus 0^*$ for the finite mask
$D_L(n), d_j$	largest parallel depth on a slice; the rank dynamic program	$R = \{(i, m-1-i) : i < m\}$ Part I's $d_L(w)$ (depth of one word), its D in Theorem 9.4 and $D[j]$ in Proposition 6.1
Γ	the alphabet $\{0, 1, @ \}$	the same as Part I's Γ ; Part III's graph Γ_R

Positions inside a word are zero based here ($v = v_0 \cdots v_{\ell-1}$), as in Part III; Part I indexes letters from one and boundaries from zero, and the interval notation $w[i : j]$ is the same in all three Parts.

14 The question and the contribution

14.1 A precise gap in the repository

[write, batch 66] “The ProveIt report” of this section is Part I.

The ProveIt report *Aligned Fragments and Constrained Crossover* (Part I) develops a fragment-rank calculus for constrained crossover. Its regular-input section (Section 7) constructs a distance automaton with at most $s4^s$ states (Theorem 7.3) and obtains an EXPSPACE upper bound for finite stabilization from a complete s -state DFA (Corollary 7.4). Its Section 11 question headed *Sharp regular-input complexity* (Question 11.2) asks whether the special two-sided subset structure permits a better algorithm and asks for matching hardness results for DFA and NFA inputs.

This is the question pursued here. The inspected repository revision is

b8b0fa2184a044d46ce9ed0f25f88d7bb60fa042.

The relevant report is at

SetTheory/Cardinals/docs/reports/
automata-and-formal-languages/constrained-crossover-closure/.

The source and README were inspected, with the regular-input conclusion and its explicit open question retrieved separately. We do not present the older fragment formulas, the coordinate-hull observation, or the report’s context-free undecidability theorem as new results.

Hughes’s preprint [1] supplies the surrounding language-operation setting. The present problem concerns *self*-crossover with equal parent lengths and identical cut positions. It is not unrestricted splicing, insertion degree, or arbitrary-length crossover. In particular, we do not claim to be solving anew the context-free one-step question already answered in the repository report (Theorem 4.2).

14.2 Main theorem

An NFA may have several initial states and no ε -transitions. A single-initial-state convention gives the same complexity results, as explained in Section 20. Let $s \geq 1$ be its number of states. Set

$$m = s + 2, \quad T_s = 2m^2 + 2m, \quad p_s = \text{lcm}(1, 2, \dots, m), \quad B_s = 2T_s + 1 = (2s + 5)^2. \quad (27)$$

The constants are deliberately conservative. The relevant qualitative facts are a quadratic transient, a period with polynomially many bits, and a rank bound independent of that period.

Theorem 14.1 (Main result). *Let L be recognized by an s -state NFA.*

1. *Finite stabilization of parallel constrained self-crossover is decidable in PSPACE. It is PSPACE-complete already for a fixed binary alphabet, even for factorial languages with universal coordinate hull. The same statements hold for frozen-source self-crossover.*
2. *If stabilization occurs, every word of the full closure has aligned fragment rank at most B_s . Hence the parallel sequence stabilizes by generation $\lceil \log_2 B_s \rceil$ and the frozen-source sequence by generation $B_s - 1$.*
3. *If stabilization does not occur, there are constants $c > 0$, C , and a residue r modulo p_s such that, for every sufficiently large $n \equiv r \pmod{p_s}$, the maximum finite rank on length n lies between $cn - C$ and n .*
4. *For DFA input, a periodic safe-core criterion decides the same property. It gives an interior obstruction of length at most $p_s s^2$ whenever stabilization fails. The exact complexity for DFA input is not determined here.*

The proof is divided into independent components. Section 15 recalls and proves the rank interpretation. Sections 16–18 establish the structural dichotomy. Section 19 derives the space upper bound, and Section 20 proves binary hardness. Section 21 handles the deterministic refinement.

14.3 What changes, and what does not

Question	Inspected report	This article
DFA finite stabilization	Decidable; EXPSpace upper bound	PSPACE upper bound; periodic safe-core criterion
NFA finite stabilization	Sharp complexity asked explicitly	PSPACE-complete over two letters
Bound in the stabilizing case	Reduction to generic limitedness	Rank at most $(2s + 5)^2$; logarithmic generation bound
Nonstabilizing regular inputs	Unbounded finite rank	Linear maximum rank on a live residue class
Exact DFA complexity	Open in the report	Still not settled here

[Added 1 October 2026, batch 73: The last row, and item (4) of Theorem 14.1, which leaves the exact DFA complexity undetermined, are out of date. Part VI settles the exact complexity for DFA input: finite stabilization of an explicit complete DFA is coNP-complete, already over a binary alphabet (Theorem 67.1).]

The article’s mathematical claims are accompanied by proofs rather than an assertion that no equivalent theorem has ever appeared. The bounded source audit covers the pinned report, Hughes, the unary-periodicity input, and the NFA universality input. It does not exclude all equivalent statements in recombination, symbolic dynamics, or semigroup literature. The software checks finite instances; it neither proves the complexity lower bound nor substitutes for the universal arguments.

15 Crossover, aligned fragments, and the two clocks

15.1 Definitions, including the empty word

Fix a finite nonempty alphabet Σ . For words x, y of a common length n and a boundary $0 \leq i \leq n$, an aligned crossover output is $x[0 : i]y[i : n]$. Both choices of parent order are allowed, and endpoint cuts are permitted. For languages A, B define

$$A \otimes_c B = \{x[0 : i]y[i : n], y[0 : i]x[i : n] : x \in A_n, y \in B_n, 0 \leq i \leq n\},$$

where the displayed set ranges over all $n \geq 0$ and $A_n = A \cap \Sigma^n$. Define

$$T(L) = L \otimes_c L, \quad L^{[0]} = L, \quad L^{[k+1]} = T(L^{[k]}),$$

and the frozen-source sequence

$$Y_0 = L, \quad Y_{k+1} = Y_k \otimes_c L.$$

Both sequences preserve the set of lengths present in L ; endpoint cuts then make both sequences increasing. These are different clocks. In a parallel generation both parents may already be recombinants; in a frozen generation one parent must be a seed.

Definition 15.1 (Aligned fragment rank). For a nonempty word w of length n , a nonempty interval $[i, j]$ is *admissible* if some $v \in L_n$ satisfies $v[i : j] = w[i : j]$. The rank $\rho_L(w)$ is the smallest number of admissible intervals in a partition

$$0 = i_0 < i_1 < \cdots < i_d = n.$$

Set $\rho_L(w) = \infty$ if no such partition exists. For the empty word, set $\rho_L(\varepsilon) = 1$ when $\varepsilon \in L$, and ∞ otherwise.

The total length and the coordinates in this definition are indispensable. An ordinary factor occurring somewhere in a seed of an unrelated length need not be an admissible aligned fragment.

For $n \geq 1$, the coordinate hull is

$$\text{Rec}(L)_n = \prod_{i=0}^{n-1} \{a \in \Sigma : \exists v \in L_n (v_i = a)\}. \quad (28)$$

If L_n is empty the product is empty. At length zero we stipulate $\text{Rec}(L)_0 = L_0$, rather than using the unconditional empty-product convention. Thus ε is never introduced from an empty length-zero seed.

[write, batch 66] The operation, the two clocks, Definition 15.1 and (28) restate Part I’s Definition 2.1, Theorem 3.8’s iteration, Definition 3.1 (“admissible” here is Part I’s “ L -compatible”) and Definition 2.4, with letters indexed from zero. They are printed because the proofs below are written in this form; they define nothing new.

15.2 Rank formulas, credited background

The following formulas are part of the earlier ProveIt report, Part I (Theorems 3.2, 3.3 and 3.8, and Corollary 3.5). Their proofs are included to make the dependency on rank completely transparent. [write, batch 66] The statements are printed once, in Part I; the proposition and corollary below restate them in Part II's notation, and their proofs are a second route, by the manuscript. Part III proves the hull formula (31) by a third route (Lemma 31.3).

Proposition 15.2 (The fragment interpretation; Part I, restated). *For every $k \geq 0$,*

$$L^{[k]} = \{w : \rho_L(w) \leq 2^k\}, \quad Y_k = \{w : \rho_L(w) \leq k + 1\}. \quad (29)$$

Moreover,

$$\rho_{T(L)}(w) = \left\lceil \frac{\rho_L(w)}{2} \right\rceil \quad (30)$$

with the natural convention for infinity, and

$$\bigcup_{k \geq 0} L^{[k]} = \bigcup_{k \geq 0} Y_k = \text{Rec}(L). \quad (31)$$

Second proof (manuscript 01). A fragment of a word in $T(L)$ meets at most two pieces inherited from seed parents. Intersect it with that parent's crossover cut, discarding any empty intersection. Therefore a partition into d $T(L)$ -fragments gives one into at most $2d$ seed fragments.

Conversely, pair consecutive pieces of a shortest seed-fragment partition. For a pair, choose a seed parent supplying each piece and cross the parents at their shared boundary. The pair is then one admissible $T(L)$ -fragment. A final unpaired piece remains admissible because $L \subseteq T(L)$. This proves (30). Iteration proves the first formula of (29), since rank at most one is exactly seed membership.

A frozen crossover appends, or prepends, at most one seed fragment to an inherited rank partition, so every word in Y_k has rank at most $k + 1$. For the reverse inclusion, take a partition into $d \leq k + 1$ seed fragments and corresponding parents $v_1, \dots, v_d$. Starting with v_1 , cross successively with $v_2, \dots, v_d$ at the partition boundaries. After the j th parent has been used, the prefix through that parent's fragment is correct; the last parent supplies the remaining suffix. This constructs w in $Y_{d-1} \subseteq Y_k$.

Finally, a nonempty word has finite rank exactly when each of its coordinates is supplied by a length-matched seed: singleton pieces give the reverse implication. That is (28). The empty word follows from its stipulated convention. $\square$

Corollary 15.3 (Finite stabilization is bounded rank). *Put*

$$R(L) = \sup(\{1\} \cup \{\rho_L(w) : w \in \text{Rec}(L)\}).$$

The parallel sequence stabilizes at a finite generation if and only if $R(L) < \infty$, and the same is true for the frozen-source sequence. When finite, the least stabilization indices are respectively $\lceil \log_2 R(L) \rceil$ and $R(L) - 1$.

Second proof (manuscript 01); Part I's Corollary 3.5 and Theorem 3.8. Both operators are monotone and extensive along their respective sequences. An adjacent equality therefore makes all later stages equal, so a stabilized value must be the union (31). The rank thresholds (29) give all the assertions, including $L = \emptyset$ through the extra element 1 in the supremum. $\square$

The trivial slice bound is $\rho_L(w) \leq |w|$ for nonempty hull words. What is not trivial is a bound uniform in word length. The rest of the article explains why regularity forces such a bound to be small whenever it exists.

16 The exact-length geometry of an NFA

Let $\mathcal{A} = (Q, \Sigma, (\Delta_a)_{a \in \Sigma}, I, F)$ be an ε -free NFA, with $|Q| = s$. A transition relation acts on subsets by

$$\Delta_a(X) = \{q' : \exists q \in X (q, q') \in \Delta_a\}.$$

For a word x , let Δ_x be the composed relation, with Δ_ε the identity. Put $E = \bigcup_a \Delta_a$. Forgetting labels but counting edges defines the exact-length reachability sets

$$P_i = E^i(I), \quad R_j = \{q : \exists f \in F (q, f) \in E^j\}. \quad (32)$$

Thus $P_0 = I$, $R_0 = F$, $P_{i+1} = E(P_i)$, and $R_{j+1} = E^{-1}(R_j)$.

For a fixed total length n , the states which can occur after i letters in some accepting length- n computation are

$$S_{n,i} = P_i \cap R_{n-i} \quad (0 \leq i \leq n). \quad (33)$$

An NFA run may be chosen separately for each admissible fragment. There is no requirement that different fragments use compatible seed runs.

Lemma 16.1 (Aligned fragment criterion). *For $0 \leq i < j \leq n$, a word x of length $j - i$ can occur as an aligned seed fragment on $[i, j]$ if and only if*

$$\Delta_x(P_i) \cap R_{n-j} \neq \emptyset. \quad (34)$$

Equivalently, there is a path labelled x through the successive viable layers $S_{n,i}, \dots, S_{n,j}$.

Proof. An accepting seed run supplies the path and hence (34). Conversely, a path counted by (34) extends backwards by some length- i path from an initial state and forwards by some length- $(n - j)$ path to a final state. The labels of those extensions may be arbitrary. The resulting accepting word has the prescribed fragment at precisely the prescribed coordinates. Every intermediate state of this completed run lies in its viable layer, which proves the equivalent formulation. $\square$

Corollary 16.2 (Coordinate alphabets). *The letters available at position i of a length- n seed are exactly*

$$A_{n,i} = \{a \in \Sigma : \Delta_a(S_{n,i}) \cap S_{n,i+1} \neq \emptyset\}. \quad (35)$$

For $n \geq 1$, $\text{Rec}(L)_n = \prod_{i < n} A_{n,i}$.

Only two unary processes, P_i and R_j , encode the dependence on the absolute length. The labelled transitions still encode arbitrary nondeterminism, but away from the two boundaries their allowed state sets will be periodic.

16.1 A direct fixed-word rank algorithm

For later audits, it is useful to record the elementary dynamic program. Given $w \in \Sigma^n$, precompute P_i and R_j for $i, j \leq n$ and set $d_0 = 0$. For $j = 1, \dots, n$, let

$$d_j = 1 + \min\{d_i : 0 \leq i < j, \Delta_{w[i:j]}(P_i) \cap R_{n-j} \neq \emptyset\}, \quad (36)$$

where a minimum over the empty set is infinity. Then $d_n = \rho_L(w)$ for $n > 0$. Maintaining the subset image as j increases for each fixed i implements the interval tests in polynomial time. The finite verifier independently compares this path-based algorithm with literal seed-substring enumeration; they do not share an interval oracle. [Added 1 October 2026, batch 73: Part V proves that the longest admissible prefix may always be taken (Lemma 57.1), which turns this program into a left-to-right pass in $O((|\Sigma| + n)s^2)$ time (Corollary 58.2).]

17 Uniform periodicity and its quantitative cost

17.1 Elementary eventual periodicity

The recurrence taking (P_i, R_i) to (P_{i+1}, R_{i+1}) is deterministic on at most 4^s pairs. Thus there exist $t \geq 0$ and $p \geq 1$ such that

$$P_{i+p} = P_i, \quad R_{i+p} = R_i \quad (i \geq t), \quad (37)$$

and one can choose $t + p \leq 4^s$. This elementary observation suffices for the qualitative structural theorem and, with polynomial-bit counters, for the space upper bound. A stronger classical unary result supplies the quadratic rank bound.

17.2 The imported unary theorem, with the normalization exposed

We use the corrected unary arithmetic-progression theorem of To [8]. In its normalized m -state form, the accepted lengths are a finite union of singleton exceptions no larger than $2m^2 + m$ and arithmetic progressions whose positive periods are at most m and whose initial terms are smaller than $2m^2 + 2m$. These constants are taken from the *corrected* proof in Section 3 of that paper, not the earlier argument whose gap the paper identifies.

Lemma 17.1 (A common quadratic transient). *For an s -state NFA, (37) holds with $t = T_s$ and $p = p_s$ as defined in (27). Moreover $\log_2 p_s = O(s \log(s + 2))$.*

Proof. Fix a target state q . Regard E as a unary NFA transition relation. Add a new source with one unary edge to each state of I , and a new sink with one unary edge from q . The new source has no incoming edges and the new sink no outgoing edges. A path of length i from I to q in the old graph exists exactly when the normalized graph accepts length $i + 2$. There are $m = s + 2$ states. Empty path languages cause no difficulty.

Beyond T_s , every singleton exception has disappeared and every positive progression is active. Membership is unchanged on adding a common multiple of all its periods. Since every period is at most m , p_s is such a common multiple. The harmless length shift by two is already past the same threshold whenever $i \geq T_s$. Thus membership of q in P_i is p_s -periodic there. The same normalization with source q and target set F proves the assertion for R_i . Taking all states does not change the common threshold or the common period.

Finally $p_s \leq m! \leq m^m$, giving the bit-length estimate. No optimal estimate for the least common multiple is needed. $\square$

The imported theorem is used only for this quantitative improvement. The structural mechanism below uses arbitrary t, p satisfying (37); it does not require a minimal period, nor does a large period force a large stabilization rank.

18 Periodic interiors and the obstruction dichotomy

Fix t, p satisfying (37). All phase arithmetic in this section is modulo p . Define

$$P_a^\infty = P_i \quad (i \geq t, i \equiv a \pmod{p}), \quad R_a^\infty = R_i \quad (i \geq t, i \equiv a \pmod{p}).$$

These definitions are independent of the representative. A residue r is *live* when

$$P_r^\infty \cap F \neq \emptyset. \quad (38)$$

Let $\mathcal{R}$ denote the set of live residues. For $n \geq t$, $L_n \neq \emptyset$ precisely when $n \bmod p \in \mathcal{R}$.

18.1 The finite periodic system

Fix a live residue r . At phase a , put

$$S_{r,a} = P_a^\infty \cap R_{r-a}^\infty, \quad (39)$$

$$C_{r,a} = \{c \in \Sigma : \Delta_c(S_{r,a}) \cap S_{r,a+1} \neq \emptyset\}. \quad (40)$$

The set $S_{r,a}$ is the viable state set at a sufficiently interior boundary of phase a in total length residue r . The set $C_{r,a}$ is the corresponding coordinate alphabet.

Lemma 18.1 (Nonvacuity and exact agreement). *For live r , all sets $S_{r,a}$ and $C_{r,a}$ are nonempty. If $n \equiv r \pmod{p}$ and $t \leq i \leq n - t$, then $S_{n,i} = S_{r,i \bmod p}$. When both boundaries of a letter lie in that interior, $A_{n,i} = C_{r,i \bmod p}$.*

Proof. The equalities follow directly from periodicity and (33)–(35). For nonemptiness, choose an arbitrarily large accepted length in residue r and an interior boundary of phase a . An accepting run supplies a viable state and, with one further interior position available, a permitted letter. $\square$

Definition 18.2 (Phase-admissible and liftable words). A word $v = v_0 \cdots v_{\ell-1}$ is *phase-admissible* at (r, a) if $v_j \in C_{r,a+j}$ for every j . It is *liftable* if there is a path with label v through $S_{r,a}, S_{r,a+1}, \dots, S_{r,a+\ell}$, starting at any state of $S_{r,a}$. A phase-admissible word that is not liftable is an *interior obstruction*.

Equivalently, start with $Z_0 = S_{r,a}$ and compute

$$Z_{j+1} = \Delta_{v_j}(Z_j) \cap S_{r,a+j+1}. \quad (41)$$

Liftability is $Z_\ell \neq \emptyset$. Every phase-admissible single letter is liftable by definition of $C_{r,a}$, so an obstruction has length at least two. This small fact is important when cuts are counted later.

18.2 Universal interiors imply a bounded number of fragments

Theorem 18.3 (Interior universality criterion). *For a regular seed language L , the following are equivalent:*


1. *Parallel self-crossover stabilizes after finitely many generations.*
2. *Frozen-source self-crossover stabilizes after finitely many generations.*
3. *At every live (r, a) , every phase-admissible word is liftable.*

When these conditions hold,

$$R(L) \leq 2t + 1. \quad (42)$$

Proof of sufficiency and the rank bound. Let $w \in \text{Rec}(L)$, with $n = |w|$. The empty word has rank one. If $1 \leq n \leq 2t$, singleton fragments give rank at most $2t$. Suppose $n > 2t$. Then $r = n \bmod p$ is live. Every letter of $w[t : n - t]$ belongs to the corresponding $C_{r,a}$ by Lemma 18.1, so this central word is phase-admissible. By condition (3), it is liftable. Lemma 16.1 makes the entire central word one admissible seed fragment.

Partition the first t and last t positions into singletons, retaining the central fragment as one piece. This gives at most $2t + 1$ pieces. If $t = 0$, the central word is the entire nonempty word and the same argument gives rank one. Corollary 15.3 proves both forms of stabilization. The necessity is established in the next theorem. $\square$



18.3 One obstruction forces linear rank growth

Theorem 18.4 (Periodic pumping of an obstruction). *Suppose v is an interior obstruction at a live pair (r, a) . There are fixed words u, b, z with $b \neq \varepsilon$ and $|b| \equiv 0 \pmod{p}$ such that, for every $k \geq 1$,*

$$w_k = ub^kz \in \text{Rec}(L), \quad \rho_L(w_k) \geq k + 1. \quad (43)$$

More strongly, for every sufficiently large $n \equiv r \pmod{p}$ there is $w \in \text{Rec}(L)_n$ with

$$\rho_L(w) \geq \left\lfloor \frac{n - |u| - |z|}{|b|} \right\rfloor + 1. \quad (44)$$

Proof. Extend v on the right by arbitrary phase-admissible letters until its length is a positive multiple h of p . This is possible because every coordinate alphabet is nonempty. Call the result b . It starts at phase a and returns to phase a , so b^k is phase-admissible for every k . Its first $|v|$ positions in each copy form the same obstruction.

Choose lengths $U, V \geq \max(1, t)$ such that

$$U \equiv a \pmod{p}, \quad U + V \equiv r \pmod{p}.$$

The length $U + V$ is at least t and has live residue, so choose an accepted word $uz \in L_{U+V}$ with $|u| = U$, $|z| = V$.

We first check all coordinates of ub^kz . For a prefix position $i < U$, the forward length i has not changed. Its suffix length has increased by kh , which is a multiple of p , and even the shorter original suffix length is at least $V \geq t$. Thus (34) for that single letter has the same truth value as in the accepted word uz . For suffix letters the argument is symmetric: the suffix length is unchanged and the prefix length increases by kh , starting at least $U \geq t$. All middle letters are in the stable alphabets $C_{r,a+j}$. Hence every coordinate is allowed, and $ub^kz \in \text{Rec}(L)$.

Now consider any partition of ub^kz into admissible seed fragments. If one occurrence of v had no cut strictly between its first and last letters, that whole occurrence would be contained in one fragment. Restricting the fragment's witnessing seed run to this occurrence would lift v through its stable viable layers. That contradicts the choice of v . Therefore every occurrence contains an internal partition cut. The k occurrences are disjoint and each has length at least two, so their internal boundary sets are disjoint. There are at least k cuts, hence at least $k + 1$ fragments. This proves (43) and, with Corollary 15.3, completes the necessity in Theorem 18.3.

For an arbitrary sufficiently large $n \equiv r \pmod{p}$, write

$$n - U - V = kh + d, \quad 0 \leq d < h.$$

Both $n - U - V$ and h are multiples of p , so d is a multiple of p . After b^k , insert any phase-admissible word e_d of length d . It also returns to phase a . The same coordinate argument proves $ub^ke_dz \in \text{Rec}(L)_n$, and the same k disjoint occurrences give (44). $\square$

Corollary 18.5 (Residue-wise bounded/linear dichotomy). *For $n \geq 1$, write*

$$M_L(n) = \max\{\rho_L(w) : w \in \text{Rec}(L)_n\},$$

using $M_L(n) = 0$ for an empty slice. On each live residue r , exactly one of the following alternatives holds:

1. $M_L(n) \leq 2t + 1$ for all $n > 2t$ in that residue;

2. $M_L(n) = \Theta(n)$ as $n \rightarrow \infty$ in that residue.

The second alternative occurs precisely when that residue has an interior obstruction. It can occur in some residues and not others.

Proof. Apply the central-fragment proof separately to each residue with no obstruction. Otherwise combine (44) with $M_L(n) \leq n$. These alternatives cannot overlap. $\square$

[Added 1 October 2026, batch 73: Part V sharpens alternative (2) to $M_L(n) = \sigma_r n + O(1)$ with a rational slope $\sigma_r > 0$ (Theorem 62.1).]

Corollary 18.6 (Asymmetric boundary bound). *If P_i is p -periodic for $i \geq t_P$ and R_i is p -periodic for $i \geq t_R$, then finite stabilization implies*

$$R(L) \leq t_P + t_R + 1.$$

Proof. Use the same stable state sets with separate valid representatives. The pumping proof works with $U \geq \max(1, t_P)$ and $V \geq \max(1, t_R)$. If no obstruction exists, keep the central interval $[t_P, n - t_R)$ as one fragment and use singleton boundary pieces. Shorter words require at most $t_P + t_R$ singletons. $\square$

19 Polynomial-space decision and finite obstruction bounds

19.1 A finite graph, represented without expansion

Fix live r and a starting phase a . The subset evolution (41) is a directed graph on configurations

$$(j, Z), \quad j \in \mathbb{Z}/p\mathbb{Z}, \quad Z \subseteq Q.$$

From (j, Z) a letter $c \in C_{r,j}$ gives

$$(j, Z) \longrightarrow (j + 1, \Delta_c(Z) \cap S_{r,j+1}). \quad (45)$$

The initial configuration is $(a, S_{r,a})$, and an obstruction is exactly a path to an empty-subset configuration.

Lemma 19.1 (A short finite obstruction). *If an interior obstruction exists at (r, a) , one exists of length at most $p2^s$. It may be padded to a multiple-of- p block of length at most $p2^s$.*

Proof. A shortest path to an empty subset repeats no nonterminal configuration: a repeated phase and subset could be deleted without changing the remainder of the path, including its phase-dependent alphabet constraints. There are at most $p2^s$ configurations. Rounding the resulting positive length up to a multiple of p does not exceed $p2^s$, since the latter is itself a multiple of p . $\square$

This bound controls an *interior obstruction*, not a rejected word of the original language and not the length of a full crossover witness. The exterior words in Theorem 18.4 are still needed.

19.2 The PSPACE upper bound

Theorem 19.2 (Polynomial-space upper bound). *For an explicitly given NFA over an explicitly given finite alphabet, finite constrained self-crossover stabilization, under either clock, belongs to PSPACE.*

Proof. Choose the canonical $t = T_s$ and $p = p_s$. Store p in binary. It has $O(s \log(s + 2))$ bits. For a residue a , a suitable representative is

$$i_a = t + ((a - t) \bmod p).$$

The sets P_a^∞ and R_a^∞ are computed from E^{i_a} by Boolean matrix powering. The exponent has polynomially many bits and the matrices have s^2 entries, so this requires polynomial space (and polynomial time in the bit length of the exponent).

To recognize *nonstabilization*, nondeterministically select residues $r, a < p$. Check that r is live using (38). Initialize $j = a$ and $Z = S_{r,a}$. Repeatedly choose a letter c , verify $c \in C_{r,j}$, and update (j, Z) by (45). Accept if Z becomes empty. Terminate and reject that branch after $p2^s$ steps. The counter, phase, residue, subset, transition data, and matrix-powering workspace all have polynomial size in the input.

Lemma 19.1 and Theorem 18.3 show that this nondeterministic machine accepts exactly the nonstabilizing inputs. Thus nonstabilization is in NPSPACE. Savitch's theorem [10] gives $\text{NPSPACE} = \text{PSPACE}$, and deterministic polynomial space is closed under complement. Stabilization is consequently in PSPACE too. Corollary 15.3 identifies the two clocks. $\square$

For completeness, the resource conversion may also be seen directly. Reachability in an exponentially large graph with polynomial-bit vertices can be decided by recursively guessing a midpoint and splitting a path length in half. A deterministic search enumerates the midpoint encodings. The recursion depth is logarithmic in the number of configurations, hence polynomial, and each level stores polynomially many bits. This is the configuration-graph form of the cited theorem, not an assumption that nondeterministic time and deterministic time are interchangeable.

19.3 Quadratic rank and a finite generation cutoff

Corollary 19.3 (Uniform finite horizon). *Let $K_s = \lceil \log_2 B_s \rceil$. For an s -state NFA language, the following are equivalent:*

$$\text{finite parallel stabilization,} \quad L^{[K_s]} = \text{Rec}(L), \quad L^{[K_s]} = L^{[K_s+1]}.$$

For the frozen-source sequence the corresponding cutoff is $B_s - 1$.

Proof. Lemma 17.1 and Theorem 18.3 give $R(L) \leq B_s$ in the stabilizing case. Use (29). Conversely, any adjacent equality stabilizes the sequence by monotonicity, independently of the chosen cutoff. $\square$

This is a *uniform cutoff*, not a proposal to construct the enormous intermediate automata for all generations. The polynomial-space algorithm tests the periodic interior directly. It is the structural criterion, not an unrestricted inclusion test on exponentially expanded automata, that keeps the space bound polynomial.

Corollary 19.4 (An explicit full-witness length bound). *If an s -state NFA does not stabilize, then, for every integer $q \geq 1$, there is a hull word of rank at least $q + 1$ and length at most*

$$2T_s + 2p_s - 2 + qp_s 2^s.$$

Proof. Use Lemma 19.1 for a block of length at most $p_s 2^s$. In the pumping construction, choose U and V in $[T_s, T_s + p_s - 1]$ with their required residues. The word $ub^q z$ has the asserted length and rank. $\square$

In particular, rank just beyond the polynomial bound has a witness of singly exponential length, with a succinct periodic description. The word need not itself have a polynomial-length explicit representation.

20 Binary PSPACE-hardness

20.1 A factorial-language amplification lemma

A language is *factorial* if every contiguous factor of every word in the language also belongs to the language. The language of an NFA in which every state is both initial and final is factorial: restrict any accepting path to the interval carrying the desired factor.

Lemma 20.1 (Missing factors cannot be repaired in bounded rank). *Let $L \subseteq \{0, 1\}^*$ be factorial and contain $0^* \cup 1^*$. Then $\text{Rec}(L) = \{0, 1\}^*$. Moreover, finite crossover stabilization is equivalent to $L = \{0, 1\}^*$.*

Proof. For every n and every coordinate, 0^n and 1^n supply both possible letters, and $\varepsilon \in L$. Hence the hull is universal.

If L is not universal, choose $v \notin L$. Because L contains both unary languages, v has length at least two. No word of L can contain v as a factor. For the target v^k , each of its k disjoint displayed copies of v must therefore contain an internal partition cut: a seed fragment containing an entire copy would force $v \in L$. Consequently $\rho_L(v^k) \geq k + 1$. The target is in the universal hull, so rank is unbounded. Corollary 15.3 rules out stabilization. The universal language is already fixed, proving the converse. $\square$

20.2 The reduction over two letters

We use the published result of Kao, Rampersad, and Shallit [9, Lemma 6]: nonuniversality of binary NFAs with every state both initial and final is PSPACE-complete. By complement closure, universality for that class is PSPACE-complete as well. That hardness theorem is an imported complexity result, not reproved by finite testing here.

Theorem 20.2 (Binary hardness with a universal hull). *Finite crossover stabilization is PSPACE-hard on binary NFAs. Hardness persists when all states are initial and final, the language is factorial, and its coordinate hull is universal. Both iteration clocks have this property.*

Proof. Let $\mathcal{A}$ be a binary all-initial/all-final NFA recognizing K . Form the disjoint union with one state having only a 0-loop and one state having only a 1-loop. Declare both added states initial and final. The new automaton recognizes

$$L = K \cup 0^* \cup 1^*. \quad (46)$$

It adds exactly two states. Its language is factorial, and its hull is universal by Lemma 20.1.

If K is universal, so is L . Conversely, if $x \notin K$, then $x01 \notin K$, since otherwise factoriality would put its factor x in K . The word $x01$ contains both letters, so it is in neither 0^* nor 1^* . Thus L is nonuniversal. We have proved

$$K = \{0, 1\}^* \iff L = \{0, 1\}^* \iff L \text{ stabilizes.}$$

The last equivalence is Lemma 20.1; the reduction is linear in the automaton size. This proves the result for both clocks. $\square$

Corollary 20.3 (Exact complexity and least alphabet). *Finite constrained self-crossover stabilization is PSPACE-complete for NFA input over a fixed binary alphabet. Over a unary alphabet every language, regular or not, is already fixed by constrained self-crossover. Thus two letters is the least alphabet size permitting the hardness result.*

Proof. Combine Theorems 19.2 and 20.2. Over a unary alphabet there is at most one word of each length, so two same-length parents, if present, are identical. No crossover adds a word. $\square$

The reduction also handles the standard single-initial-state NFA encoding. Add one fresh initial state, make it accepting because $\varepsilon \in L$, and on each letter give it the union of the transitions out of the old initial states. Leave every old transition in place and make only the new state initial. This ε -free construction adds one state and preserves the language. The all-initial/all-final promise itself uses the multiple-start convention, but the ordinary PSPACE-completeness statement does not depend on that convention.

20.3 A second reduction connecting to the repository amplifier

The preceding proof is the sharper alphabet result. The following four-letter version explains directly how the old report's delimiter amplification also yields a hardness proof from ordinary binary NFA universality. Its block mechanism is credited to that earlier line of argument; it is not a second claimed discovery of the amplifier. [write, batch 66] The amplifier is Part I's Theorem 5.1, with the context-free languages G_K and B_K of (6)–(7); the regular languages $\widehat{G}(K)$, $\widehat{B}(K)$ below (the manuscript's $G(K)$, $B(K)$) are different: $\widehat{G}(K)$ adds to K every word containing @, with no 00 guard, and $\widehat{B}(K)$ leaves its first and last blocks unrestricted.

Let $K \subseteq \{0, 1\}^*$ be given by an NFA, let $\Gamma = \{0, 1, @, \#\}$, and put

$$\widehat{G}(K) = K \cup \Gamma^* @ \Gamma^*, \quad \widehat{B}(K) = \Gamma^* \cup \Gamma^* \# (\widehat{G}(K) \#)^* \Gamma^*.$$

An internal complete #-delimited block must lie in $\widehat{G}(K)$; the first and last incomplete blocks are unrestricted. Standard union, concatenation, and star constructions give an NFA of size linear in the input, and ε -elimination preserves the state set.

Proposition 20.4 (Exact delimiter ranks). *The hull of $\widehat{B}(K)$ is $\{0, 1, @, \#\}^*$. If K is universal, so is $\widehat{B}(K)$. If $x \notin K$, then for every $q \geq 1$,*

$$\rho_{\widehat{B}(K)}(\#(x\#)^q) = q + 1.$$

Proof. A word with at most one delimiter belongs to $\widehat{B}(K)$. For any prescribed letter at any prescribed position, fill every other position with @ to obtain such a seed. This proves the hull assertion, including the empty word.

If K is universal, $\widehat{G}(K) = \Gamma^*$ and the displayed expression for $\widehat{B}(K)$ permits every word. If $x \notin K$, each displayed $\#x\#$ in $\#(x\#)^q$ is forbidden inside a seed. Consecutive occurrences may share a delimiter letter, but their *internal boundary slots* are disjoint. Hence every fragment partition needs at least q cuts. For the opposite bound, partition as q copies of $\#x$ and one final $\#$. Each piece contains at most one delimiter and extends to a length-matched seed by filling outside positions with @. These are $q + 1$ admissible fragments. If $x = \varepsilon$, the same argument partitions $\#^{q+1}$ into single letters and remains valid. $\square$

This establishes the four-letter reduction without the all-initial/all-final hardness theorem. It is included for comparison and proof portability; the binary theorem above remains the stronger final result.

21 Deterministic automata and safe cores

Assume now that each $\Delta_c(q)$ has at most one state. The original DFA may be complete, but after restriction to viable phase sets some transitions become undefined. Fix a live residue r and regard

$$(a, q), \quad q \in S_{r,a},$$

as vertices of a partial deterministic phase graph. Its allowed input letters at phase a are $C_{r,a}$.

Definition 21.1 (Safe core). A family $(U_a)_{a \bmod p}$ with $U_a \subseteq S_{r,a}$ is safe if

$$q \in U_a, c \in C_{r,a} \implies \delta(q, c) \text{ is defined and belongs to } U_{a+1}.$$

The safe core $(\mathcal{U}_{r,a})$ is the greatest such family.

It is obtained by starting with all viable vertices and repeatedly deleting any vertex with an allowed letter leading outside the surviving set. The process terminates after at most ps vertex deletions. An efficient implementation stores reverse edges and a queue of newly unsafe vertices.

Theorem 21.2 (Safe-core characterization). *For a deterministic automaton and a fixed live residue r , all phase-admissible words are liftable at every starting phase if and only if the safe core is nonempty. A nonempty safe core meets every phase.*

Proof. If a safe vertex exists, any allowed letter takes it to a safe vertex in the next phase. Following one allowed letter at a time, possible because all $C_{r,a}$ are nonempty, shows that safe vertices exist in every phase. Starting with any safe vertex at a desired phase then reads every phase-admissible word without leaving the viable graph. This proves sufficiency.

Suppose the safe core is empty. A vertex is safe exactly when no phase-admissible continuation can cause its deterministic run to leave the viable graph. Equivalently, every vertex now has some killing continuation. One can see this directly from deletion times: a vertex removed first has an immediately invalid allowed transition; a later removed vertex has a transition to an earlier removed vertex.

Start with all viable states at some phase. Select one survivor and append a killing word for that survivor. The statewise determinism implies that every current state has at most one resulting image. Since the selected state has no image, the number of surviving images drops by at least one. Repeat with a surviving state at the new phase. After at most s rounds there are no survivors. The concatenated word is phase-admissible but not liftable. Thus universality fails. $\square$

Corollary 21.3 (A shorter deterministic obstruction). *If a live residue of a deterministic s -state automaton has an interior obstruction, it has one of length at most ps^2 .*

Proof. When the safe core is empty, a shortest route from any given vertex to an invalid allowed transition traverses no vertex twice before failure. It has length at most the number of phase vertices, at most ps . Use such a route in each of at most s killing rounds in the proof of Theorem 21.2. $\square$

For a fixed residue with its phase sets explicitly listed, the safe core can be computed in $O(ps|\Sigma|)$ time by backward deletion. Treating all p residues separately gives $O(p^2s|\Sigma|)$ core work, in addition to constructing the phase sets. These are bounds in the *expanded periodic presentation*. Since p can be exponentially large relative to the original state count, this is not a polynomial-time algorithm for arbitrary DFA input. The PSPACE upper bound still follows from Theorem 19.2; no matching DFA lower bound is claimed.

[Added 1 October 2026, batch 73: Part VI proves the problem coNP-complete for explicit complete DFAs (Theorem 67.1), with a polynomial certificate of nonstabilization built from SCC stripes rather than from the expanded phase graph (Proposition 70.1).]

Example 21.4 (Why the deterministic argument cannot be used on NFAs). Take two states x, y , both initial and final, with transitions

$$\Delta_0(x) = \{x, y\}, \quad \Delta_1(x) = \emptyset, \quad \Delta_0(y) = \emptyset, \quad \Delta_1(y) = \{x, y\}.$$

The language is all binary words: after any letter the full two-state subset is available again. Yet neither individual state supports both possible first letters. A vertexwise safety deletion removes both states, even though the subset evolution never empties. The deterministic cardinality-drop argument fails because a single NFA state can branch into several surviving states. This is a concrete reason to retain subset universality in the NFA algorithm.

22 Quantitative consequences and worked examples

22.1 The generation bound has the right logarithmic order

Theorem 22.1 (Logarithmic generation lower bound). *For each $m \geq 1$, let*

$$L_m = \{0^{i-1}10^{m-i} : 1 \leq i \leq m\}.$$

This language has a complete DFA with at most $2m + 3$ states, stabilizes finitely, and satisfies $R(L_m) = m$. Its least parallel stabilization index is $\lceil \log_2 m \rceil$, and its least frozen-source index is $m - 1$.

Proof. A DFA stores the current length $0, \dots, m$, whether a 1 has already occurred, and one rejecting sink. Moving past length m or seeing a second 1 enters the sink. This uses $2(m + 1) + 1$ states; unreachable states may be removed but are irrelevant to the bound.

All nonempty hull words have length m , so rank is at most m . Every coordinate of 1^m is supplied by one of the seed words. No admissible fragment of 1^m can have length at least two, because a seed has only one 1. Hence $\rho_{L_m}(1^m) = m$. The case $m = 1$ is included, with rank one. Apply Corollary 15.3. $\square$

Thus the largest possible finite parallel stabilization index for s -state automata has order $\Theta(\log s)$, even over two letters. This does *not* prove that the quadratic rank bound is optimal. The presently proved lower and upper orders for maximal finite rank are linear and quadratic, respectively. [Added 1 October 2026, batch 71: Part IV raises the lower order to quadratic for NFAs: Theorem 47.1 gives an s -state NFA with finite maximal rank $(s - 3)^2$, so the quadratic rank bound has the optimal order (its constant is open), and the parallel index of order $\Theta(\log s)$ is $2 \log_2 s + O(1)$ (Section 52). The family L_m above has a complete DFA; for DFAs the lower order in this report stays linear (Question 53.2).] [Added 1 October 2026, batch 73: Part VI gives complete binary DFAs with $m + 6$ states and finite maximal rank at least $(m - 1)^2 + 2$ (Theorem 67.2), so the lower order for DFAs is quadratic as well, with leading constant one.] [Added 1 October 2026, batch 73: The constant left open above is one (Part VII, Theorem 76.2).]

22.2 A regular language with a repairable boundary

Example 22.2 (One generation without universality). Over $\Sigma = \{0, 1\}$, let

$$L = \{\varepsilon\} \cup 0\Sigma^* \cup \Sigma^*0.$$

Then $\text{Rec}(L)_1 = \{0\}$ and $\text{Rec}(L)_n = \Sigma^n$ for $n \geq 2$. For any $w \in \Sigma^n$ with $n \geq 2$, choose an interior cut. The prefix is supplied by a seed ending in 0 and the suffix by a seed starting in 0. Hence rank is at most two and $T(L) = \text{Rec}(L)$. The word 11 is not a seed and has rank exactly two, so the least parallel and frozen-source indices both equal one.

This example separates finite stabilization from universality for general regular languages. The hardness equivalence with universality applies to the special factorial class, not to all NFAs.

22.3 Bounded and linear residues can coexist

Example 22.3 (Even lengths are fixed; odd lengths have linear rank). Let

$$L = \{w \in \{0, 1\}^* : |w| \text{ is even}\} \cup 0^* \cup 1^*.$$

Its hull is universal. At even lengths every word is a seed, so the maximum rank is one. At odd lengths the seeds are exactly $0^n, 1^n$; an admissible nonempty fragment must be constant. Therefore the rank of an odd-length word is its number of maximal constant runs, and

$$M_L(n) = \begin{cases} 1, & n \text{ even,} \\ n, & n \text{ odd.} \end{cases}$$

Here $M_L(0)$ may also be taken as one. A four-state NFA recognizes the language by taking a disjoint union of a parity automaton and two unary loops. The periodic criterion sees a universal interior at the even total residue and an obstruction such as 01 at the odd residue.

The example shows why a statement of linear growth at *all* large lengths would be false. It also explains why the total-length residue r and the boundary phase a are separate parameters in the construction.

22.4 The strongest possible generation growth in a bad residue

On a live residue with an obstruction, let

$$D_L(n) = \max_{w \in \text{Rec}(L)_n} \min\{k : w \in L^{[k]}\}.$$

By (29) and Corollary 18.5,

$$D_L(n) = \log_2 n + O(1) \quad (n \rightarrow \infty \text{ in that residue}). \quad (47)$$

The maximum frozen-source depth is $M_L(n) - 1 = \Theta(n)$. The constants can be read from the pumped block length and boundary lengths. In a good residue the two depths remain bounded. Thus the rank dichotomy yields a complete bounded-versus-maximal-order dichotomy for the two natural iteration clocks, residue by residue.

23 Executable audits and their limits

The package contains a standard-library Python implementation in `code/verify.py`. It is an exact finite reference checker, not an optimized prover and not a formalization. Its NFA transitions and subsets are represented by integer bit masks. The recorded run used Python 3.13.5 and returned PASS; details and limitations are stored in `data/verification.json` and the console transcript `data/verification.txt`.

[write, batch 66] *Shipped names*. In this report the delivered files are shipped, byte-identical, as follows; the manuscript and its PDF are printed or replaced by this report, and the delivery README by the report's README. The unprefixd data/verification.json of this report is Part I's receipt (Section 10).

Shipped path	Delivered as
code/02-pspace-stabilization-verify.py	code/verify.py
code/02-pspace-stabilization-example.py	code/example.py
code/02-pspace-stabilization-build.sh	build.sh
data/02-pspace-stabilization-verification.json	data/verification.json
data/02-pspace-stabilization-verification.txt	data/verification.txt
data/02-pspace-stabilization-build-receipt.json	data/build-receipt.json
02-pspace-stabilization-proof-audit.md	notes/proof-audit.md
02-pspace-stabilization-source-ledger.md	notes/source-ledger.md

23.1 Two independent rank calculations

The path-based algorithm implements (36). The independent algorithm enumerates all accepted seed words at a given length, collects their actual aligned substrings, and runs a partition dynamic program on those substring sets. In particular, it does not use P_i , R_j , or viable layers to decide interval admissibility. All tested finite ranks, including infinity and the empty-word case, agreed.

The literal crossover check enumerates all parent pairs and cuts. It then compares its actual generated sets with the rank thresholds 2^k . These checks are logically separate from the periodic criterion.

23.2 Recorded coverage

Exact audit	Recorded count
All labelled two-state binary NFAs, all initial/final subsets	4,096
Independent rank comparisons, all binary words of lengths 0–5	258,048
Stable-rank bound checks within those comparisons	69,488
Complete three-state binary DFAs, fixed initial state	5,832
Safe-core versus subset-criterion comparisons on those DFAs	5,832
Stable three-state DFA rank checks at length 6	369,152
Additional partial-deterministic core comparisons	1,296
Pumped-witness checks across the stated families	1,641
Binary all-initial/all-final hardness-transform cases	256
All binary finite seed sets on slices of lengths 0–3	278
Literal parallel-generation comparisons on those seeds	1,112
One-1 lower-family cases, $1 \leq m \leq 24$	24
Unary graph/subset cases on 1–3 vertices	4,164
Unary periodicity equalities within those cases	4,382
Exact delimiter ranks, including the omitted empty word	504
Delimiter coordinate-supply witnesses	9,072
Parity-example rank checks, lengths 1–10	2,046
Boundary-repair example rank checks, lengths 0–7	255

The two-state NFA family has 3,720 stabilizing and 376 nonstabilizing members according to the implemented exact periodic criterion. The three-state complete DFA family has 5,768 stabilizing and 64 nonstabilizing members. These are labelled-automaton counts, not counts of distinct languages or isomorphism classes.

An additional 200 seeded random NFAs with three through six states were tested, with 4,000 sampled rank queries. The seed is 20260930. Those samples are explicitly not exhaustive. The worked examples are also checked directly: the delimiter family uses every omitted binary word of length at most five and one through eight repeated blocks; the parity example uses every binary word of lengths one through ten; and the boundary-repair example uses every binary word of length at most seven. The unary audit enumerates every directed graph on one through three labelled vertices and every chosen initial subset, taking the same subset as final; it tests both forward and backward sequences through their exact eventual cycles.

The categories overlap: for example, some of the rank-bound checks are also independent rank comparisons. Their counts must not be added and reported as a number of statistically independent experiments.

23.3 What the program actually certifies

The periodic solver explicitly enumerates the least eventual cycle of (P_i, R_i) and then searches phase/subset graphs by breadth first search. On the tested small automata this is convenient, but it stores the expanded graph. It is *not* an implementation with the polynomial-space resource guarantee of Section 19. The latter is an on-the-fly theoretical algorithm with binary counters and matrix powers.

For each detected small obstruction, the program constructs the exterior accepted word and the padded repeated block from the proof, then verifies finite ranks of the resulting full words. On stabilizing cases it checks the stronger observed-transient bound $2t + 1$, where t is obtained from the least pair cycle, rather than only the much looser universal B_s .

No finite test proves that every rank above every state count is bounded, or that every obstruction pumps. Those are theorems proved in Sections 18–19. No solver is used to infer PSPACE-hardness, no verification output is represented as a Lean kernel certificate, and no literature theorem is declared proved because small instances pass.

23.4 Reproduction

From the package root, run

```
python3 code/verify.py --output data/local-verification.json
latexmk -pdf -interaction=nonstopmode -halt-on-error article.tex
```

The explicit output option preserves the delivered receipt. The default output path is `data/verification.json`; running without the option replaces that file in the local copy. There are no network calls, third-party Python dependencies, or external theorem provers in the checker. The included build script performs only the local LaTeX build.

[write, batch 66] *In this report, do not run the shipped files in place.* The verifier’s default output path is `../data/verification.json` relative to its own directory, which here is Part I’s receipt; the example program imports the verifier under its delivered module name `verify`, which the prefixed file does not have; and the build script runs `latexmk` on an `article.tex` in its own directory, `code/`, where there is none. Copy the two programs into a scratch directory under their delivered names `code/verify.py` and `code/example.py`, and run them there, with an explicit `--output`; the report’s README gives the commands and the result of such a rerun. The article is rebuilt from the report directory as in Section 10. The delivered example program’s comment calls the parity language “article Example 9.3”; in this report it is Example 22.3.

24 Further research questions

The questions below are specific continuation targets. They are not all asserted to be new or globally open. The first is the unsolved portion of the inspected report’s regular-input complexity question; the others are natural sharpenings exposed by the present proofs. [write, batch 66] The inspected report’s question is Part I’s Question 11.2. Question 24.6 below is Part I’s Question 11.1 and is printed as a pointer; Question 24.7 overlaps Part I’s Question 11.7, asking it for the dichotomy proved here.

Question 24.1 (Exact complexity for DFA input). What is the exact complexity of deciding finite stabilization for complete DFAs over a fixed binary alphabet?

The present PSPACE upper bound is stronger than the earlier EXPSPACE bound, but the NFA hardness proof does not transfer to a DFA without an exponential determinization. The safe-core characterization removes subset branching once the phases are explicit. The remaining issue is whether the exponentially long unary phase structure can be manipulated more economically, or can encode a matching hard problem. An algorithm polynomial in p does not answer this question. [Added 1 October 2026, batch 73: Answered in Part VI (Theorem 67.1): coNP-complete under polynomial-time many-one reductions, for both clocks; nonstabilization is NP-complete, and the coNP upper bound holds over every explicit finite alphabet. The exponentially long phase structure is avoided by naming SCC stripes, of which there are at most s (Theorem 69.4); the hardness encodes 3SAT through Chinese-remainder length testers (Lemma 71.1).]

Question 24.2 (The optimal finite rank bound). Let $F(s)$ be the largest finite $R(L)$ among binary languages recognized by NFAs with at most s states. Is $F(s)$ linear, quadratic, or of an intermediate order?

There are only finitely many labelled NFAs of a fixed state bound, so this maximum is well defined. For $s \geq 5$, the results here give

$$\left\lfloor \frac{s-3}{2} \right\rfloor \leq F(s) \leq (2s+5)^2.$$

A quadratic lower bound would have to exploit long unary transients while making many boundary positions require genuinely separate aligned fragments. A long transient alone is not such a proof. Conversely, a linear upper bound would need a way to compress boundary fragments that is absent from the singleton argument. [Added 1 October 2026, batch 71: Answered as to order, for NFAs, in Part IV: $(s-3)^2 \leq F(s) \leq (2s+5)^2$ for $s \geq 4$, so $F(s) = \Theta(s^2)$ (Theorem 47.1). The construction takes the route sketched above: a long unary length transient (a two-cycle length gate) together with two monochromatic components, which make every change of letter at a missing length a boundary between separate fragments. The upper bound is the one above. The leading constant and the exact $F(s)$ stay open (Question 53.1), and so does the order for DFAs (Question 53.2).] [Added 1 October 2026, batch 73: Part V lowers the upper bound to $2(s-1)^2 + 3$ (Corollary 64.2), so that $(s-3)^2 \leq F(s) \leq 2(s-1)^2 + 3$ and the leading constant lies in $[1, 2]$.] [Added 1 October 2026, batch 73: For complete DFAs Part VI determines the order and the constant: $(s-7)^2 + 2 \leq F_{\text{DFA}}(s) \leq s^2 + 3$ for $s \geq 8$ (Theorem 67.2).] [Added 1 October 2026, batch 73: Part VII proves that the leading constant is one: $(s-3)^2 \leq F(s) \leq s^2 + 8s$ for $s \geq 4$ (Theorem 76.2); $F(1), \dots, F(4) = 1, 2, 5, 7$ (Theorem 84.1).]

Question 24.3 (Exact residue-wise rank slope). On a bad residue, does $M_L(n)/n$ always have a limit, perhaps after refining the period? Must a resulting slope be rational, and can it be computed effectively?

The pumping theorem gives positive lower and upper linear bounds, not a limit or a formula. The quantity is a maximum over target words of a minimum partition cost. Replacing it by a shortest-path mean without justifying this order of quantifiers would be invalid. The periodic systems provide finite objects on which this problem can be studied. [Added 1 October 2026, batch 73: Answered in all three parts by Part V

(Theorem 62.1): on every live residue r of any valid period, $M_L(n) = \sigma_r n + O(1)$, where σ_r is rational with denominator at most $p(2^s - 1)$, is computed as the largest reachable cycle mean of a finite subset-reset graph, and is positive exactly when the residue has an interior obstruction; no refinement of the period is needed. The order of quantifiers is handled as asked above: longest-prefix greedy factorization is optimal (Lemma 57.1), so the minimum partition cost is computed by a deterministic parser before any maximization.]

Question 24.4 (Optimal obstruction lengths). How sharp are the interior bounds $p2^s$ for NFAs and ps^2 for DFAs? Can the multiplicative p be replaced by a smaller invariant of the unary reachability graph?

The NFA bound is a configuration-count bound; the DFA bound concatenates at most s single-state killing words. Neither construction is claimed optimal. The definition of an interior obstruction must be retained: shortest nonmembers of L can instead reflect only a repairable boundary.

Question 24.5 (Compressed periods and effective certificates). Can the live residues and their safe or obstructed status be represented by polynomial-size arithmetic data, rather than by all residues modulo an explicit least common multiple?

The common period used here is convenient but often highly redundant. Strongly connected component periods and arithmetic-progressions representations are candidate data structures. A positive certificate would need to establish universality in every live phase system, not merely exhibit a few successful runs. A negative certificate can already be represented by one forbidden interior word and periodic exterior padding, but that word may itself be exponentially long. [Added 1 October 2026, batch 73: For DFAs Part VI gives polynomial-size data of this kind in one direction: the stripes (H, h) of Section 69 have polynomial-size names, and nonstabilization has a polynomial certificate (Proposition 70.1). A polynomial certificate of stabilization for all DFAs would give $\text{NP} = \text{coNP}$, by Theorem 67.1. This is the write's observation; the manuscript of Part VI does not discuss this question.]

Question 24.6 (Context-free finite stabilization). [write, batch 66] This is Part I's Question 11.1, printed there: what is the exact arithmetical complexity of finite stabilization for context-free grammar input under the same aligned operation? The manuscript's comment follows.

The inspected report proves undecidability and gives an upper placement in Σ_2^0 (Theorem 5.2 and Proposition 5.4), but does not claim completeness at that level. Nothing in the regular-input analysis resolves this question: unary periodicity of the two exact-length state processes uses finiteness of the automaton state set. A reduction showing unbounded rank from the failure of a Σ_2^0 property would require a different argument.

Question 24.7 (Restricted cuts and other recombination rules). Which parts of the dichotomy survive when only a regular set of cut positions is permitted, or when a single operation exchanges several alternating intervals between two parents?

For unrestricted aligned one-point self-crossover, every singleton is available in the eventual hull. Restricted cuts may destroy that product structure. In a multi-cut operation, several pieces still come from only two parents; they are not automatically independent fragments. Thus one cannot simply replace the factor two in the rank law by the number of cuts plus one without proving an appropriate realization theorem.

Question 24.8 (Proof-assistant implementation). Can the periodic-interior equivalence, quantitative cutoff, and binary hardness reduction be formalized with an auditable separation between combinatorial proofs and imported complexity theorems?

The development proposed in the next section gives small independent milestones. Formalizing the full PSPACE lower bound includes the published all-initial/all-final universality reduction, not merely the two-state union construction used here. [Added 1 October 2026, batch 71: Part IV’s Question 53.4 asks the same for its lower-bound construction.]

25 A formalization and implementation roadmap

No Lean or Rocq files are claimed as verified in this delivery. A practical formal development can nevertheless be organized without importing the entire earlier report as an unexamined axiom.

Finite words and partitions. Define an aligned interval by two bounded natural numbers and a proof of their strict order. Define fragment admissibility by a length-matched witness word. The two rank thresholds can first be stated as existence of a partition with at most d pieces, avoiding extended-natural minima until the basic equivalences have been proved. The empty-word convention should be isolated at the start.

Runs and exact-length reachability. Represent an NFA by finite relations on a finite state type. Prove Lemma 16.1 by concatenating and restricting finite runs. This is the bridge at which absolute coordinates and total lengths must be retained. The periodic theorem can initially take t, p and the periodicity equations as explicit hypotheses.

The structural core. Prove the central-fragment sufficiency result first. Then prove the pumping lemma in three separate pieces: phase-aligned block extension, preservation of each exterior coordinate under a multiple-of-period length change, and the injection from obstruction copies to partition cuts. The last piece uses the strict internal boundary sets; it does not need a vague claim that copies are “independent”.

Quantitative periodicity. The elementary 4^s eventual-cycle argument is a short finite-state lemma and can yield a fully formal qualitative result early. The quadratic bound should be added only after a verified unary arithmetic-progression theorem is available, or else recorded as an explicit conditional corollary. The two added normalization states and the length shift by two must not disappear from the statement.

Decision complexity. A reference explicit algorithm is much easier to validate than its space bound. Separate functional correctness of the phase/subset search from the representation-size proof for binary counters, Boolean matrix powering, and nondeterministic configuration graphs. Likewise, prove the binary language equivalence in (46) independently of importing PSPACE-completeness of the source problem.

Deterministic safety. The greatest safe family is a finite greatest fixed point. Its elementary delete-to-failure proof should be particularly tractable. The proof that killing one deterministic start decreases the surviving-subset cardinality should be a named lemma, making the exact point where nondeterminism breaks the argument explicit.

These milestones produce useful formal theorems before a full complexity library is available. They also keep the distinction between an existing repository formalization, an ordinary mathematical proof, and a future formalization plan visible.

26 Conclusion

The special two-sided structure identified in the repository question really does improve on generic limitedness. Exact-length reachability becomes periodic from both ends. In a good interior, all allowed target letters can be joined by one seed run, leaving only a bounded exterior cost. In a bad interior, one finite failure repeats and forces a fresh partition cut each time. This yields both the bounded/linear dichotomy and a finite search problem on phase/subset configurations.

As a consequence, finite stabilization for NFA input is PSPACE-complete over the smallest nontrivial alphabet, while every stabilizing s -state instance has rank at most $(2s + 5)^2$. The logarithmic order of the parallel generation bound is sharp. The exact DFA complexity, optimal rank order, and precise slopes of bad residues remain substantive next problems, rather than being hidden inside an unspecified generic limitedness bound. [Added 1 October 2026, batch 71: The optimal rank order is now known for NFAs: it is quadratic (Part IV, Theorem 47.1); its constant and the DFA case remain.] [Added 1 October 2026, batch 73: The slopes of bad residues are now determined: they are rational and effectively computable (Part V, Theorem 62.1), and the constant of the rank order is at most two (Corollary 64.2).] [Added 1 October 2026, batch 73: The exact DFA complexity is coNP-complete, and the rank order for complete DFAs is quadratic with constant one (Part VI, Theorems 67.1 and 67.2).] [Added 1 October 2026, batch 73: For NFAs the constant is one (Part VII, Theorem 76.2).]

27 Proof and edge-case audit

[write, batch 66] This section and the next were the manuscript's Appendices A and B.

Potential failure	How the proof handles it
Empty language or empty slice	There is no live contribution; hull slices remain empty. $R(L)$ includes 1 only to normalize the empty-language stabilization index.
The empty word	Its rank is 1 exactly when it is a seed. The length-zero hull is explicitly L_0 , not an unconditional empty product.
Transient $t = 0$	The central-fragment argument then gives rank 1 in the stabilizing case. Pumping uses positive exterior lengths to avoid vacuous baselines.
An obstruction of length one	Impossible by the definition of $C_{r,a}$; every permitted letter has a viable transition. Thus every obstruction has an internal cut slot.
A forbidden word only at a particular phase	The padded block has length divisible by p , so every copy starts at the same phase.
The total-length residue changes	It does not: all inserted block and remainder lengths are multiples of p .
Exterior coordinates cease to be available	The unchanged side of the single-letter path test is exact; the changed side lies beyond its transient and changes by a multiple of p .
A fragment's parent run is not globally fixed	Each admissible fragment has its own seed witness. Restricting that witness is sufficient to contradict an interior obstruction.
Linear growth only on a sparse subsequence	The remainder-padding step supplies every sufficiently large length of the chosen residue, not merely lengths of the form $U + kh + V$.
Confusing a huge period with a huge rank bound	The period is used to build and recognize the interior. Rank cost comes from the two transients, which are quadratic independently of the period.
Replacing NFA universality by statewise safety	The two-state example in Section 21 disproves that replacement. Safety is used only for deterministic transitions.

Potential failure	How the proof handles it
Binary hardness filled by unary loops	A missing word x in a factorial language forces the missing mixed word $x01$, so adding $0^* \cup 1^*$ cannot create false universal cases.
Sharing a delimiter in the second reduction	The forbidden occurrences share a letter, not an internal boundary slot. Their required cuts remain distinct, including $x = \varepsilon$.
Claiming a polynomial-time DFA algorithm	The safe-core time is polynomial only in the expanded period. The article does not identify it with polynomial time in the original automaton.
Finite checks presented as a proof	They are bounded audits. Universal claims rest on the written arguments and the explicitly cited imported theorems.

28 Source ledger and novelty boundary

Repository source. The pinned source is the crossover report’s `article.tex` and `README.md` at the revision displayed in Section 14. The EXPSPACE upper bound is its corollary headed *Decidable regular-input stabilization*; the continuation target is its Section 11 question headed *Sharp regular-input complexity*. The rank formulas and frozen-source calculus are credited background. The report itself states that it is unrefereed and not formalized. [write, batch 66] These are Part I’s Corollary 7.4 and Question 11.2; the source ledger is shipped as `02-ospace-stabilization-source-ledger.md`.

Primary external inputs. The aligned operation was checked against Hughes’s version 1 preprint. The corrected unary constants were checked in To’s Section 3, including the page images of the correction. The binary all-initial/all-final universality hardness was checked against Lemma 6 of the version 2 Kao–Rampersad–Shallit preprint. Savitch’s space simulation is the standard complexity conversion used for the upper bound; the argument’s configuration-graph application is explained in the text.

What is imported, and what is proved here. The unary arithmetic-progression theorem and the PSPACE-hardness of the source universality problem are imported published results. The periodic-interior criterion, obstruction pumping theorem, rank and depth consequences, crossover complexity reduction, and deterministic safe-core application are proved in this article. The earlier rank calculus is reproved for completeness, not claimed as a new contribution.

Search boundary. Searches on 30 September 2026 included combinations of constrained crossover, finite stabilization, regular languages, recombination, PSPACE, unary arithmetic progressions, and all-initial/all-final NFA universality. Broad recombination searches were noisy and did not amount to an exhaustive literature classification. The defensible novelty claim is a proved extension of the inspected repository, not an unconditional assertion of worldwide priority or a claim of peer-reviewed acceptance.

No hidden formal status. No assertion in this article was checked by Lean or Rocq. The Python program and its receipts are included so that the finite audits can be reproduced, but the ordinary proofs remain independently reviewable. No repository file was modified as part of this delivery.

[write, batch 66] The last sentence describes the delivery; the write phase added this Part and dated notes in Part I. Placement beside ProveIt’s Lean and Rocq developments confers no formal status on Part II,

and none of its statements is formalized. No ProveIt Lean or Rocq file defines finite automata, crossover, aligned rank, unary arithmetic-progression normal forms or space complexity classes, so To's theorem, Kao–Rampersad–Shallit's Lemma 6 and Savitch's theorem enter only as cited mathematics.

Part III. Decidable regularity and exact stabilization of crossover closures

[write, batch 66] *Part III was added on 30 September 2026 in batch 66, together with Part II, from a second manuscript written against Part I as printed above (it did not know Part II). For context-free seeds it decides regularity, inclusion and equality of the full crossover closures; for semilinear one-marker seeds it computes every generation, decides finite stabilization, and classifies the stabilizing ones into regular, linear non-regular and non-context-free closures; and it shows exact stabilization depth to be a clique number, coNP-complete to threshold for succinct ray input. It partly answers Part I’s Question 11.4. It shares with Part II only a reproof of Part I’s coordinate hull, printed here as a third route (Lemma 31.3). Section 29 records the provenance and fixes the notation of the addition.*

29 [write, batch 66] Provenance and notation of Part III

29.1 One manuscript, one addition

Part III is batch-66 manuscript 03, *Decidable Regularity and Exact Stabilization of Crossover Closures: arithmetic invariants, sparse-marker trichotomies, and immediate loss of context-freeness*, dated 30 September 2026 (23 US-letter pages as delivered). It is an AI-assisted research manuscript: its title page and PDF metadata name ChatGPT as the drafting assistant and Vladimir Reshetnikov as the person it was prepared for. It arrived as the archive `ProveIt_Crossover_Classification.zip` (inner directory of the same name) in ProveIt commit 9ad899cbe and was placed in commit 4fee1cd07. Like Part II it was written against ProveIt commit b8b0fa218 (b8b0fa2184a044d46ce9ed0f25f88d7bb60fa042), at which Part I’s `article.tex` had the blob eb4f427f printed above, so “the preceding report” of this Part is Part I. Its code, recorded results, ray examples and audit notes are shipped beside Part I’s files with the prefix `03-closure-classification-`; its manuscript, delivery README and PDF are not shipped and survive in the arrival commit, and its checksum list SHA256SUMS (21 files, all verified at placement) was retired by the repository’s policy. Section 45 and the report’s README list the shipped names.

Every result, proof, example, remark, table, limitation and question of the manuscript is printed here, in its order. Its title page, abstract and status box became this section; its Sections 1–14 are Sections 30–43; its two appendices became Sections 44 and 45. Its nine unlabelled research questions received labels. Text marked [write, batch 66] was added when the manuscript was written into this report: the notes relating it to Part I and Part II, among them the credit to Part I’s Lemma 8.1 after Lemma 31.4, which the manuscript does not give. No theorem was strengthened or weakened.

29.2 The manuscript’s summary and status

Abstract, as delivered (“the ProveIt report on this operation” is Part I). Constrained crossover exchanges suffixes of equal-length words at an aligned cut. The ProveIt report on this operation proves that the infinite self-crossover closure is a product of position supports. We use that invariant to decide regularity, inclusion, equality, and eventual equality of the full closures of arbitrary context-free seeds. Regularity is exactly finite row type of the associated Presburger support relations; a negative answer has an effective infinite family of distinguishable marked prefixes.

For the linear context-free seeds $L_R = 0^* \cup \{0^i 10^j : (i, j) \in R\}$, where R is semilinear, we determine

every generation and characterize finite stabilization. In the finitely stabilizing case, we obtain a complete regular/linear context-free/non-context-free trichotomy. Two co-occurring, distinct interior marker velocities cause non-context-freeness already after one crossover. For finite unions of arithmetic rays the exact maximum marker count is a clique number, giving coNP-completeness of stabilization by an input depth for binary-encoded ray descriptions, even when the seed and closure are regular. An explicit binary linear seed stabilizes after one step at a non-context-free language with at most four words of each length. Complete proofs, executable exact ray analysis, 12,867 finite checks, and further research questions are included.

Status and attribution, as delivered. Unrefereed research manuscript, not a Lean or Rocq development. General monadic decomposition and the slender-language background are classical ingredients, not claimed discoveries. The crossover applications and sharp sparse-marker results are proposed contributions relative to the inspected repository; global priority is not established.

29.3 Notation beside Part I and Part II

Part III uses Part I's operation $\otimes_c$, its parallel iterates $L^{[g]}$ and its frozen-source iterates $Y_g = L \otimes_c Y_{g-1}$. Positions are zero based. Seven symbols were renamed when the manuscript was written in; no normalization changed:

- the full closure $\mathcal{H}(L) = \bigcup_g L^{[g]}$ is printed as Part I's $\text{Rec}(L)$, which is the same set by Theorem 3.3 (and $\mathcal{H}_R$ as Rec_R);
- the crossover map $C(L) = L \otimes_c L$ is printed as Part I's $T(L)$;
- the gap-ray language $T_s(a, d)$ of (66) is printed as $\text{Gap}_s(a, d)$, and the finite union T of Lemma 36.1 as $\mathcal{T}$, because T is the crossover map and Part II's T_s is a constant;
- the row equivalence $E_a(i, r)$ of (55) is printed as $\text{Eq}_a(i, r)$, because Part I's E_a is the position support, here called P_a^L ;
- the family $L_k = L_{R_k}$ of (75) is printed as L_k^{fam} , because Part II's L_m is another family;
- the threshold $T = 2^d + 1$ in the proof of Theorem 39.2 is printed as θ .

In Part III	Meaning in Part III	Not to be confused with
$P_a^L(i, j)$	some seed has a with i letters before and j after (48)	the same relation as $(i, j) \in E_a$ in Part I (11); not Part I's coordinate set $P_{n,i}(L)$ or Part II's state set P_i
$A_L(n, i)$	letters allowed at zero-based position i	Part I's $P_{n,i+1}(L)$ (one-based); Part II's $A_{n,i}$, the same set
$\text{Eq}_a, U_{a,r}, V_{a,r}, C_a$	row equivalence, its classes and rows, least representatives (Theorem 33.1)	Part I's E_a, U_i, V_i ; Part II's $C_{r,a}$
$R \subseteq \mathbb{N}^2$	a marker mask	Part I's $R_L(n)$; Part II's $R(L)$ and R_j
$K_R(n), k_R(n), K(R)$	marker sites at length n , their number, its supremum	Part I's languages K, K_q ; Part II's cutoff K_s
$d_{\parallel}(R), d_{\text{frozen}}(R)$	stabilization depths of L_R	Part I's depth $d_L(w)$ of one word
$\text{Gap}_s(a, d), \nu(d)$	gap-ray language with s ones, number of growing gaps	Part II's constant T_s
Γ_R, ω	ray compatibility graph, clique number	the alphabet $\Gamma = \{0, 1, @\}$ of Parts I and II
E, F	exceptional lengths, finite points of R	Part II's edge relation E and final states F
M_a	the marked language of Theorem 33.2	Part II's $M_L(n)$
$B_{k,g}, L_k^{\text{fam}}$	slice counts and the family of Theorem 40.2	Part I's B_K ; Part II's B_s and L_m (which is L_R for a finite R)
S_i	DFA states reachable by length- i words (proof of Theorem 33.1)	Part I's S_q ; Part II's $S_{n,i}$

30 Research target and relation to ProveIt

30.1 The question being pursued

A context-free language can have a non-context-free constrained-crossover closure, while the length generating function of that closure remains rational. Thus enumeration alone does not identify its language-theoretic complexity. The natural next question is whether the special product structure of the closure makes its regularity decidable, despite the undecidability of regularity for arbitrary context-free languages. A second question is quantitative: when the closure is reached in finitely many generations, what is the exact depth, and how does it interact with loss of context-freeness?

We answer the first question for every context-free seed. We answer the second, including a sharp language-class classification, for all semilinear one-marker seeds with uniformly bounded position support. These are substantial families of linear context-free languages, and they include regular seeds with succinctly specified arithmetic periods.

The inspected ProveIt revision is

b8b0fa2184a044d46ce9ed0f25f88d7bb60fa042.

The directly preceding report is *Aligned Fragments and Constrained Crossover*, dated 29 September 2026, in

SetTheory/Cardinals/docs/reports/
automata-and-formal-languages/constrained-crossover-closure/.

Its source and README were inspected at that revision. Our global-regularity question is a continuation motivated by that report; we do not attribute it to a numbered open question unless explicitly stated in the source.

[write, batch 66] The preceding report is Part I. Its Question 11.4 (*Grammar classes preserving context-freeness*) asks which context-free seeds have a context-free coordinate hull and suggests restrictions on the position supports E_a ; Theorem 37.1 and Corollary 37.2 answer it completely for finitely stabilizing semilinear one-marker seeds, by exactly such a restriction, and Theorem 33.1 decides the regular case for every context-free seed. The general question stays open (Question 42.1). The manuscript’s attribution policy above is kept: it does not claim to answer a numbered question.

30.2 What is inherited, and what is developed here

The preceding report proves a coordinate-hull description, aligned-fragment rank formulas, undecidability of one-step equality for context-free seeds, undecidability of eventual stabilization, regular-input algorithms, and rational commutative counting for closures. It also gives a linear context-free seed with a non-context-free closure and unbounded stabilization depth. We do not claim those results again. We reprove the short coordinate-hull fact because every later argument depends on its exact normalization, especially empty slices and the empty word. [write, batch 66] These are Part I’s Theorems 3.2–3.8, 4.2 and 5.2, Section 7, Theorem 8.3 and Section 9. Part II, written in parallel, adds a PSPACE-complete regular-input stabilization problem; Part III does not use it.

Hughes explicitly left the single-step constrained-crossover question open in Section 4 of his August 2026 preprint [1]. That is the question the preceding ProveIt report claims to answer. The present manuscript is not a second priority claim about the same question. Instead it develops full-closure decision procedures and a finitely stabilizing classification that are not among the preceding report’s stated results.

Input / problem	Result proved here	Location
Arbitrary CFG seeds	Inclusion, equality, eventual equality, and all Boolean-combination length sets are effective	§32
Regularity of the full closure	A decidable finite-row criterion, DFA synthesis, or an infinite obstruction certificate	§33
Semilinear one-marker seed	Exact generations and finite-stabilization criterion	§34–35
Sparse semilinear support	Complete regular / linear nonregular / non-context-free trichotomy	§37
Binary-encoded arithmetic rays	Exact clique formula; NP-complete support threshold and coNP-complete depth threshold	§38–39
One-step loss of context-freeness	Binary linear seed, depth one, at most four words per length	§40

30.3 Classical ingredients and the novelty boundary

Effective semilinearity of Parikh images and its equivalence with Presburger definability are standard [11, 3, 12]. The finite-row criterion used below is a concrete instance of monadic decomposition, not a new general solution of that problem. Relevant prior work includes monadic decomposability of regular relations [13], integer-linear-arithmetic decomposition [14], and the July 2026 results on the size of decompositions [15]. In particular, decidability does not imply that a small decomposition or small DFA exists.

The structural argument for sparse languages is closely related to the classical paired-loop description of slender context-free languages; Raz gives a constructive treatment [16]. We supply a self-contained finite-ray pumping argument, then derive the exact velocity criterion and its generation consequences for crossover. CLIQUE hardness itself is classical [17]; the contribution here is the realization of an arbitrary graph by distinct fixed marker positions with arithmetic length constraints, together with the exact depth identity.

The literature search was bounded. Equivalent sparse-recombination formulations elsewhere have not been exhaustively excluded. Mathematical correctness is supported by the proofs, not by a claim of first discovery.

31 Coordinate supports and the full closure

Throughout, $\mathbb{N} = \{0, 1, 2, \dots\}$ and Σ is a finite nonempty alphabet. Positions are zero based. A letter at position i in a word of length n has i letters before it and $j = n - 1 - i$ after it.

Definition 31.1. For $A, B \subseteq \Sigma^*$, constrained crossover is

$$A \otimes_c B = \{wz, yx : wx \in A, yz \in B, |w| = |y|, |x| = |z|\}.$$

Endpoint cuts are allowed. Put $T(L) = L \otimes_c L$, $L^{[0]} = L$, $L^{[g+1]} = T(L^{[g]})$, and $\text{Rec}(L) = \bigcup_{g \geq 0} L^{[g]}$. The *parallel stabilization depth* is the least g with $L^{[g]} = L^{[g+1]}$, or infinity if no such g exists.

Since endpoint cuts are allowed, $L \subseteq T(L)$. Monotonicity shows that one adjacent equality implies all later equalities and equality with the full closure. This is the parallel clock, not the frozen-source clock $Y_0 = L$, $Y_{g+1} = L \otimes_c Y_g$.

[write, batch 66] This definition restates Part I's Definition 2.1 and Lemma 2.2; the manuscript defines the full closure as the union of the generations, which is Part I's coordinate hull $\text{Rec}(L)$ by Theorem 3.3, hence the shared symbol. The relation P_a^L defined next is Part I's position support: $P_a^L(i, j)$ holds exactly when $(i, j) \in E_a$ (11).

Definition 31.2. The genuine position support of a language is

$$P_a^L(i, j) \iff \exists u, v \in \Sigma^* : |u| = i, |v| = j, uav \in L. \quad (48)$$

For $n > 0$ let $A_L(n, i) = \{a \in \Sigma : P_a^L(i, n - 1 - i)\}$. The empty-word bit is $e_L = 1$ precisely when $\varepsilon \in L$.

Lemma 31.3 (Coordinate hull; inherited fact, reproved). *For every language L and every $n > 0$,*

$$\text{Rec}(L) \cap \Sigma^n = \{a_0 \cdots a_{n-1} : a_i \in A_L(n, i) \text{ for all } i\}. \quad (49)$$

Moreover $P_a^{\text{Rec}(L)} = P_a^L$ for every a , and $\varepsilon \in \text{Rec}(L)$ if and only if $\varepsilon \in L$.

[write, batch 66] This is Part I's Theorem 3.3 (with Definition 2.4), restated in support form; Part II restates and reproves it as (31). The manuscript's proof, printed next, is a third route: Part I contracts ranks, Part II pairs fragments, and this proof builds the target one letter at a time from seed donors, as in Part I's frozen-source Theorem 3.8.

Third proof (manuscript 03). Crossover preserves length and copies each coordinate from a parent of that same length. Thus no generation creates a new letter support, proving one inclusion in (49). For the reverse inclusion, choose for each position i a seed word whose letter there is a_i . Starting with a donor for the first position, take a prefix already matching positions $0, \dots, i - 1$ and a suffix from a donor for position i , crossing at i . This retains the correct prefix and adds one correct letter. The other parent is a seed, hence belongs to every later generation; after at most $n - 1$ such steps the desired word is obtained. If a length slice of L is empty, all its genuine supports are empty and both sides of (49) are empty. The support equality follows from $L \subseteq \text{Rec}(L)$ and preservation. Only a length-zero seed can supply a length-zero offspring. $\square$

Lemma 31.4 (Effective arithmetic support). *Given a context-free grammar for L , one can effectively compute Presburger formulas defining every P_a^L and compute e_L .*

[write, batch 66] This is Part I's Lemma 8.1, which the manuscript does not cite: the same effective semilinearity of the same supports (P_a^L is Part I's E_a), there proved through an inverse homomorphism and a regular intersection, here reproved by a transducer argument. The added clause on e_L is the classical nullable-nonterminal test.

Second proof (manuscript 03). A finite-state transducer nondeterministically marks exactly one occurrence of a , maps each preceding letter to a fresh symbol x , erases the marked letter, and maps each following letter to y . Its image is the context-free language $\{x^i y^j : P_a^L(i, j)\}$, with an effectively constructible grammar. Its Parikh image is exactly the set in (48). Effective Parikh semilinearity and the effective equivalence between semilinear sets and Presburger formulas give the claim [11, 3, 12]. The nullable-nonterminal computation decides the empty-word bit. $\square$

Remark 31.5 (Why genuine supports matter). An arbitrary assignment of masks to coordinates can specify letters at some coordinates of a dead length while leaving another coordinate empty. Those assigned letters are not realized in any word. The regularity necessity argument below uses genuine supports, not such unnormalized masks. The one-marker construction of Section 34 avoids the issue by including 0^* , so every length is live.

32 Closure comparison and Boolean length theory

Theorem 32.1 (Inclusion, equality, and eventual comparison). *Given context-free grammars for L and K , it is decidable whether $\text{Rec}(L) \subseteq \text{Rec}(K)$ and whether $\text{Rec}(L) = \text{Rec}(K)$. The sets of lengths at which inclusion or equality fails are effectively ultimately periodic. Consequently eventual inclusion, eventual equality, and finiteness of the symmetric difference are decidable.*

Proof. By Lemma 31.3, inclusion is equivalent to

$$e_L \leq e_K \quad \text{and} \quad \bigwedge_{a \in \Sigma} \forall i, j \in \mathbb{N} (P_a^L(i, j) \Rightarrow P_a^K(i, j)). \quad (50)$$

Sufficiency is coordinatewise. For necessity, if the support implication fails, a witnessing seed word of L already lies outside $\text{Rec}(K)$. This also proves that for positive lengths the failure set is exactly

$$D_{L,K} = \{n > 0 : \exists a, i, j [n = i + j + 1 \wedge P_a^L(i, j) \wedge \neg P_a^K(i, j)]\}. \quad (51)$$

The alphabet quantifier is a finite disjunction. This is an effective unary Presburger set, hence effectively ultimately periodic. Add length zero when $e_L > e_K$. Equality failure is the union of the failure sets in both directions. Presburger arithmetic decides boundedness of these sets. A language over a finite alphabet is finite exactly when its length set is bounded, so the same test decides finiteness of the symmetric difference. $\square$

Corollary 32.2 (Effective witnesses). *When inclusion fails, its least failure length is computable, and a word in $L \setminus \text{Rec}(K)$ is computable. When eventual inclusion fails, one can compute a nonempty arithmetic progression of failure lengths.*

Proof. Find the least element of the effective unary failure set; solve the support formula for a letter and two lengths, then search the finite length slice using CFG membership. A witnessing seed word must exist. An unbounded unary Presburger set contains an effectively computable infinite arithmetic progression in its eventual periodic part. $\square$

Theorem 32.3 (Boolean combinations have effective periodic length sets). *For context-free $L_1, \dots, L_m$ and any Boolean expression B in the languages $\text{Rec}(L_1), \dots, \text{Rec}(L_m)$, the set $\{|w| : w \in B\}$ is effectively ultimately periodic. In particular, emptiness, universality, and finiteness of B are decidable.*

Proof. Convert the Boolean expression into a finite disjunction of conjunctions of positive and negative membership requirements. Treat $n = 0$ by the empty-word bits. Fix one conjunction and $n > 0$. If its positive indices are I , write

$$W(a, t, n) = \bigwedge_{h \in I} P_a^{L_h}(t, n - 1 - t).$$

An empty conjunction is true. Every position must admit at least one positive-compatible letter:

$$\forall t < n \bigvee_{a \in \Sigma} W(a, t, n). \quad (52)$$

For each negative index j , choose a witness position $t_j < n$ and a letter a_j such that $W(a_j, t_j, n)$ holds and $P_{a_j}^{L_j}(t_j, n - 1 - t_j)$ fails. Require that equal witness positions carry equal chosen letters. There are only finitely many choices of the a_j ; after their finite enumeration these conditions are Presburger.

Necessity follows by reading a word meeting the conjunction. Conversely, assign the chosen letters to the witness positions. The compatibility conditions make these assignments consistent. At every remaining position choose a letter supplied by (52). The resulting word satisfies all positive and negative requirements. The existence of such a word is therefore an effective Presburger condition on n . Finite union over the conjunctions proves the theorem. Universality is emptiness of the complement; finiteness is boundedness of the length set. $\square$

These procedures compare the *full closures*; they do not decide inclusion or equality of the original context-free seeds. Nor do they solve the eventual-stabilization problem for arbitrary grammars.

33 Decidable regularity through row types

For a relation $P \subseteq \mathbb{N}^2$ let $\text{row}_P(i) = \{j : P(i, j)\}$. A relation has finite row type if only finitely many such sets occur.

Theorem 33.1 (Regularity criterion and synthesis). *For a context-free seed L , the following are equivalent:*

- (i) $\text{Rec}(L)$ is regular.
- (ii) Every P_a^L has finite row type.
- (iii) Every P_a^L is a finite union of products of unary ultimately periodic sets.
- (iv) There are $B \in \mathbb{N}$ and $p \geq 1$ such that for all a, i, j ,

$$i \geq B \implies (P_a^L(i + p, j) \leftrightarrow P_a^L(i, j)), \quad (53)$$

$$j \geq B \implies (P_a^L(i, j + p) \leftrightarrow P_a^L(i, j)). \quad (54)$$

These conditions are decidable from a grammar. In the positive case a DFA for $\text{Rec}(L)$ is effectively constructible.

Proof. Suppose a DFA with state set Q recognizes $\text{Rec}(L)$. Let S_i be the set of states reachable from the start state by arbitrary words of length i . The truth of $P_a^{\text{Rec}(L)}(i, j)$ is determined by S_i , the transition on a , and existence of an accepting continuation of length j . There are at most $2^{|Q|}$ possible sets S_i , so only finitely many rows. Since $P_a^{\text{Rec}(L)} = P_a^L$, (i) implies (ii).

Assume (ii), and define the Presburger row equivalence

$$\text{Eq}_a(i, r) \iff \forall j \in \mathbb{N} (P_a^L(i, j) \leftrightarrow P_a^L(r, j)). \quad (55)$$

Choose one representative r per class. The sets $U_{a,r} = \{i : \text{Eq}_a(i, r)\}$ and $V_{a,r} = \{j : P_a^L(r, j)\}$ are unary Presburger sets. Thus $P_a^L = \bigcup_r U_{a,r} \times V_{a,r}$, proving (iii). A common threshold and a common multiple of the finitely many unary periods prove (iv). Conversely, separate periodicity partitions each axis into the finitely many values below B and the residue classes above B ; P_a^L is constant on each product cell, proving (iii) directly from (iv).

Given (iii), refine the finitely many unary sets to finite partitions on both axes. The complement of P_a^L is also a finite union of unary ultimately periodic rectangles, say $\bigcup_h U_h \times V_h$. For a unary length set U , write $\Sigma^U = \{u : |u| \in U\}$, a regular language. The words having at least one forbidden coordinate form

$$\bigcup_{a \in \Sigma} \bigcup_h \Sigma^{U_h} a \Sigma^{V_h},$$

which is effectively regular. Its complement has exactly the positive length slices of $\text{Rec}(L)$; adjust membership of ε using e_L . This proves (i) and gives DFA synthesis.

For effectivity, form the set of least row representatives

$$C_a = \{i : \forall r < i \neg \text{Eq}_a(i, r)\}. \quad (56)$$

It is Presburger, and finite row type is exactly boundedness of C_a , expressed by $\exists B \forall i > B (i \notin C_a)$. Presburger decision therefore settles (ii). In the positive case, find a bound, enumerate the representatives, and perform the construction above. Alternatively, (iv) itself is a Presburger sentence: $i + p$ and $j + p$ use addition, not multiplication of variables. $\square$

Theorem 33.2 (Effective infinite obstruction). *If $\text{Rec}(L)$ is nonregular, one can compute a letter a and integers $b \geq 0$, $p \geq 1$ such that the rows $\text{row}_{P_a^L}(b + pm)$, $m \geq 0$, are pairwise distinct. Equivalently,*

$$\forall m \neq m' \exists j \quad P_a^L(b + pm, j) \triangle P_a^L(b + pm', j),$$

where $\triangle$ denotes exclusive disjunction. This gives infinitely many pairwise distinguishable prefixes of a marked rational image of $\text{Rec}(L)$.

Proof. Some canonical representative set C_a is infinite. Its effective unary semilinear form contains an infinite progression $b + p\mathbb{N}$, all of whose members are different representatives. Consider the marked language

$$M_a = \{0^i 10^j : P_a^L(i, j)\}.$$

The prefixes $0^{b+pm}1$ have pairwise different residual languages, as the distinguishing suffixes 0^j show. The language M_a is a rational transduction of $\text{Rec}(L)$: choose one a , relabel it 1, and relabel all other letters 0. A regular $\text{Rec}(L)$ would give regular M_a , impossible. The Presburger formulas and the progression are an effective certificate; a distinguishing j for any pair can be found by arithmetic search. $\square$

Remark 33.3 (Complexity is representation dependent). The proof uses classical monadic decomposition in a specific language application [13, 14, 15]. It gives no polynomial bound in the size of the original grammar. Effective Parikh conversion, quantifier elimination, period extraction, and explicit DFA construction may all expand the representation greatly. The polynomial claim later in this article concerns a different input format: a list of arithmetic rays.

34 One-marker seeds: exact generations

Definition 34.1. For $R \subseteq \mathbb{N}^2$, define

$$L_R = 0^* \cup \{0^i 10^j : (i, j) \in R\}, \quad (57)$$

$$K_R(n) = \{i \in \mathbb{N} : i < n, (i, n-1-i) \in R\}, \quad (58)$$

$$k_R(n) = |K_R(n)|, \quad K(R) = \sup_{n \geq 0} k_R(n). \quad (59)$$

Here $K_R(0) = \emptyset$, and $K(R)$ may be infinite. Write $\text{Rec}_R = \text{Rec}(L_R)$.

Proposition 34.2 (Every semilinear mask has a linear seed). *If R is semilinear, then L_R is effectively linear context-free.*

Proof. For a linear component $(b, c) + \mathbb{N}(u_1, v_1) + \cdots + \mathbb{N}(u_t, v_t)$ use one nonterminal A with

$$A \longrightarrow 0^{u_\ell} A 0^{v_\ell} \quad (1 \leq \ell \leq t), \quad A \longrightarrow 0^b 10^c.$$

The number of each recursive production used is arbitrary, so the exponents are exactly those of the component. Finite union of these linear grammars and a grammar for 0^* gives (57). A period $(0, 0)$ can be discarded. The construction is effective, but writing 0^u as u explicit terminals is not polynomial in a binary encoding of u . $\square$

Theorem 34.3 (Exact parallel and frozen-source generations). *For every $R \subseteq \mathbb{N}^2$, every $g \geq 0$, and every binary word $w \in \{0, 1\}^*$,*

$$w \in \text{Rec}_R \iff \text{supp}_1(w) \subseteq K_R(|w|), \quad (60)$$

$$w \in L_R^{[g]} \iff w \in \text{Rec}_R \text{ and } |w|_1 \leq 2^g, \quad (61)$$

$$w \in Y_g \iff w \in \text{Rec}_R \text{ and } |w|_1 \leq g + 1. \quad (62)$$

The empty word belongs to every language here. Hence finite stabilization occurs precisely when $K(R) < \infty$. Its exact depths are

$$d_{\parallel}(R) = \lceil \log_2 \max(1, K(R)) \rceil, \quad d_{\text{frozen}}(R) = \max(0, K(R) - 1). \quad (63)$$

Proof. Zeros are always available because $0^* \subseteq L_R$, and the allowed one positions are precisely $K_R(n)$. Lemma 31.3 gives (60). At generation zero the one-count bound is one. A crossover of two words with at most b ones has at most $2b$ ones. Conversely, for a target with at most $2b$ ones, cut between the b -th and $(b+1)$ -st one when both exist, or at an endpoint otherwise. Use as first parent the target's ones before the cut and as second parent its ones at or after the cut, filling other positions with zeros. Each parent is in the hull with at most b ones, and their crossover is the target. Induction proves (61). The analogous argument removes the last one and uses a single-marker donor for it, proving (62).

If $K(R) < \infty$, the first bounds covering all possible one counts give (63). If a bound has not reached $K(R)$, there exists a length with more available sites than that bound, and a word with exactly one more one witnesses strictness of the next generation. This also proves nonstabilization when $K(R) = \infty$. $\square$

For a nonempty target with r ones, the inherited aligned-fragment rank has the particularly simple value $\max(1, r)$: every donor fragment has at most one one, and cuts between successive ones attain that bound. Theorem 34.3 is also a direct proof that does not require assuming any unformalized rank theorem from the preceding report.

[write, batch 66] The rank value is Part I's (20) in general form. Theorem 34.3 and Corollary 34.4 contain three earlier examples: Part I's Example 3.6 ($R = \mathbb{N}^2$); Part I's seed S_q of Section 9, which is L_R for $R = \{(i, j) : j \geq (q-1)i + q - 1\} \cup \{(i, j) : i \geq (q-1)j + q - 1\}$, with $k_R(qm + r) = 2m$, so that the corollary gives (22) and (23); and Part II's family L_m (Theorem 22.1), which is L_R without the all-zero words for the finite mask $R = \{(i, m-1-i) : 0 \leq i < m\}$, where $K(R) = m$ gives the same depths $\lceil \log_2 m \rceil$ and $m-1$.

Corollary 34.4 (Exact slice enumeration). *For every $n, g \geq 0$,*

$$|\text{Rec}_R \cap \{0, 1\}^n| = 2^{k_R(n)}, \quad |L_R^{[g]} \cap \{0, 1\}^n| = \sum_{r=0}^{\min(k_R(n), 2^g)} \binom{k_R(n)}{r}.$$

In particular, Rec_R is slender, meaning uniformly bounded in the number of words per length, exactly when finite stabilization occurs.

Regularity and finite stabilization are different properties. For example, $R = \mathbb{N}^2$ gives the regular closure $\{0, 1\}^*$ but has $K(R) = \infty$. The next sections classify the finitely stabilizing case completely.

35 Sparse semilinear masks and arithmetic rays

Theorem 35.1 (Decidable sparsity and effective ray normalization). *Let $R \subseteq \mathbb{N}^2$ be semilinear. Then $K(R) < \infty$ if and only if, in any finite linear-set presentation of R , the nonzero period vectors within each component are collinear. In that case R is effectively a finite union of arithmetic rays and finitely many points:*

$$R = F \cup \bigcup_{v=1}^r ((b_v, c_v) + \mathbb{N}(u_v, v_v)), \quad u_v + v_v > 0. \quad (64)$$

Thus finite stabilization of semilinear one-marker seeds is decidable.

Proof. Suppose one component has nonparallel periods $p, q \in \mathbb{N}^2$. Let $s = p_1 + p_2 > 0$ and $t = q_1 + q_2 > 0$. For every $N \geq 0$ the $N+1$ points

$$b + (N-h)tp + hsq, \quad h = 0, \dots, N,$$

belong to that component and have the same coordinate sum $b_1 + b_2 + Nst$. They are distinct because $tp \neq sq$. Consequently $K(R)$ is unbounded.

If all nonzero periods of a component are collinear, its points lie on one line with a nonnegative nonzero direction. Such a line intersects any diagonal $i + j = n - 1$ in at most one point. A finite union of components therefore gives a uniform bound; singleton components have the same property. Since a component is contained in R , the first argument also shows that no presentation of a sparse R can contain a nonparallel period pair.

It remains to normalize a collinear component effectively. Write its periods as $a_1 d, \dots, a_t d$, where $d \in \mathbb{N}^2$ is primitive and each a_ℓ is positive. Divide the a_ℓ by their gcd g . The resulting nonnegative integer semigroup has gcd one and is cofinite. Here is an explicit effectivity argument. Choose its least positive generator a . The residues of nonnegative combinations modulo a form a subgroup of the finite cyclic group and, by gcd one, are all residues. A finite shortest-path search on residues supplies a represented integer h_r in each residue class r . Every integer at least $H = \max_r h_r$ is $h_r + ma$ with $m \geq 0$, hence represented. Enumerate represented integers below H exactly; all larger integers form a tail. Returning to the original coordinates gives one arithmetic ray of step gd plus finitely many points. A component with no nonzero periods is already a point. Finite union gives (64). $\square$

The normalization may enlarge the description. No polynomial bound is claimed when a general semilinear or Presburger description is the input. The finite points must be retained for exact depth: discarding them is harmless for language-class membership but can change the largest number of simultaneous marker sites.

Definition 35.2 (Length period and velocity). For an infinite ray indexed by v , set

$$n_v = b_v + c_v + 1, \quad q_v = u_v + v_v, \quad \lambda_v = \frac{u_v}{q_v}, \quad f_v(n) = \lambda_v n + b_v - \lambda_v n_v.$$

The ray contributes exactly one site, at position $f_v(n)$, precisely when $n \geq n_v$ and $n \equiv n_v \pmod{q_v}$. Its *velocity* is $\lambda_v \in [0, 1] \cap \mathbb{Q}$. It is *interior* if $0 < \lambda_v < 1$. Two rays are *length compatible* if

$$n_v \equiv n_w \pmod{\gcd(q_v, q_w)}. \quad (65)$$

Length compatibility says that the two rays occur together at arbitrarily large lengths. Velocity zero means a fixed distance from the left endpoint; velocity one means a fixed distance from the right endpoint. An interior site recedes linearly from both endpoints.

36 A finite-ray pumping principle

The following elementary statement isolates the point at which taking a union might otherwise invalidate a pumping argument. It is a specialized form of the paired-loop phenomenon in slender language theory [16].

For $s \geq 0$ and vectors $a, d \in \mathbb{N}^{s+1}$ with $d \neq 0$, define the *gap-ray language*

$$\text{Gap}_s(a, d) = \{0^{a_0+d_0t} 10^{a_1+d_1t} 1 \dots 10^{a_s+d_st} : t \in \mathbb{N}\}. \quad (66)$$

There are exactly s ones, and the gap vector is uniquely determined by the word. Write $v(d) = |\{h : d_h > 0\}|$.

Lemma 36.1 (Finite-ray trichotomy). *Let $\mathcal{T}$ be a finite union of languages (66), with possibly different s , together with finitely many exceptional words. Then:*

- (i) $\mathcal{T}$ is regular if and only if every ray in the union has $v(d) \leq 1$.
- (ii) $\mathcal{T}$ is context-free if and only if every ray in the union has $v(d) \leq 2$. In that case $\mathcal{T}$ is linear context-free.

These tests apply to any such presentation, including overlapping rays.

Proof. For sufficiency in (i), a ray with just one growing gap is a fixed word with a factor $(0^h)^*$ inserted at a fixed location, hence regular. For sufficiency in (ii), if two gaps grow, the words have the form

$$U 0^{pt} V 0^{qt} W, \quad t \geq 0,$$

for fixed words U, V, W and positive p, q . The grammar $A \rightarrow 0^p A 0^q \mid V$, with fixed outer words U, W , is linear. The cases of one growing gap, finite words, and finite unions preserve linearity.

For necessity, treat regularity and context-freeness simultaneously, with $h = 1$ and $h = 2$, respectively. Suppose a represented ray with $v(d) > h$ occurs in $\mathcal{T}$. Work in the fixed space of gap vectors with its number s of ones. Its affine line cannot coincide with the affine line of a ray with at most h growing coordinates. Each such low-support line therefore meets it in at most one point. Choose a word far enough along the high-support ray to avoid all these finitely many intersections, to exceed every exceptional word length, and to exceed a pumping length.

A regular pumping decomposition repeats one nonempty factor. A context-free pumping decomposition repeats two factors, at least one nonempty. If any pumped factor contains a one, pumping it sufficiently many times exceeds the globally bounded number of ones in $\mathcal{T}$, a contradiction. Thus all pumped nonempty factors consist only of zeros. Their repetition adds a nonzero gap vector $e \geq 0$ supported in at most h coordinates. The pumped words have gap vectors

$$x + me, \quad m = 0, 1, 2, \dots,$$

where x is the chosen original vector. Infinitely many of these vectors must belong to one ray of the finite presentation. Taking two distinct such vectors shows that its direction is parallel to e . That ray therefore has at most h growing coordinates, and its affine line contains x . This contradicts the choice of x .

The argument explicitly handles a finite union. It does not infer non-context-freeness merely from containment of a non-context-free sublanguage, which would be invalid. $\square$

Lemma 36.2 (Gap-ray normal form for sparse hulls). *For a finite union of arithmetic rays and finitely many points as in (64), both Rec_R and every finite generation $L_R^{[g]}$ are effectively finite unions of gap-ray languages and finitely many words.*

Proof. Let M be the least common multiple of the infinite rays' length periods, taking $M = 1$ if there are none. For each residue class modulo M , discard finitely many initial lengths so that all applicable rays have started and the order of their affine site functions has stabilized. Distinct affine functions cross at most once, so there are only finitely many order exceptions. Identical functions are merged on that residue class. Also discard the finitely many lengths carrying exceptional points.

Write the remaining lengths of this class as $n = N + Mt$, $t \geq 0$. The available distinct positions then have the form $i_j(t) = \alpha_j + \beta_j t$, with integers $\beta_j = M\lambda_j$ and a fixed order. Every subset of these positions can carry ones. If its selected positions are $i_1(t) < \dots < i_s(t)$, the zero gaps have slopes

$$(\beta_1, \beta_2 - \beta_1, \dots, \beta_s - \beta_{s-1}, M - \beta_s). \quad (67)$$

They are nonnegative: eventual ordering forces the site slopes to be nondecreasing. Their sum is $M > 0$. The constants are the nonnegative gaps at $t = 0$. This gives a gap ray. For the empty subset it gives the all-zero ray. There are only finitely many residues and subsets. A finite generation retains just the subsets of size at most 2^g , by Theorem 34.3. The discarded length slices are finite, so their words give the finite exceptions. $\square$

37 The sparse-marker language-class trichotomy

Theorem 37.1 (Velocity classification). *Let R be a finite union of arithmetic rays and finitely many points. Exactly one of the following alternatives holds.*

1. **Endpoint case.** *No infinite ray is interior. Then Rec_R and every $L_R^{[g]}$ are regular.*
2. **Single co-occurring interior velocity.** *There is an interior ray, but there is no length-compatible pair of interior rays with distinct velocities. Then Rec_R is linear context-free and nonregular, and every $L_R^{[g]}$ is linear context-free and nonregular.*
3. **Two distinct co-occurring interior velocities.** *There are length-compatible rays v, w with $0 < \lambda_v < \lambda_w < 1$. Then Rec_R is not context-free. Already $L_R^{[1]}$ is not context-free, and neither is any $L_R^{[g]}$ with $g \geq 1$.*

The seed L_R is always linear context-free. The alternative is decidable using a quadratic number of gcd and rational comparisons in the number of rays, with polynomial bit complexity for binary-encoded coordinates.

Proof. Use the gap-ray normal form of Lemma 36.2. In the endpoint case every selected site has velocity zero or one. The vector in (67) has at most one nonzero coordinate: the only possible change is from zero to M , including a change at an endpoint. Lemma 36.1 makes the hull and every generation regular.

In the second case, within any common length residue class, all interior sites have the same velocity λ . Thus the distinct site velocities are drawn from $\{0, \lambda, 1\}$. There are at most two positive successive differences in (67), so all gap rays are linear context-free. An interior ray supplies single-marker words whose two zero gaps both grow. Restrict the hull or any generation to words with exactly one one. This restricted language is a finite union of gap rays containing a two-growing-gap ray; by Lemma 36.1 it is nonregular. Regular languages are closed under intersection with the regular language 0^*10^* , so the hull and all generations are nonregular.

In the third case, the generalized Chinese remainder theorem gives an infinite common length progression for v and w . Beyond a finite threshold their positions are ordered and distinct. Select just these two sites. The resulting three gap slopes are positive multiples of

$$\lambda_v, \quad \lambda_w - \lambda_v, \quad 1 - \lambda_w.$$

Consequently the hull's exact-two-ones sublanguage has a gap-ray presentation containing a ray with three growing gaps. It is not context-free by Lemma 36.1. The same exact-two-ones sublanguage occurs in every generation $g \geq 1$, by Theorem 34.3. Intersecting a hypothetical context-free generation with the regular language $0^*10^*10^*$ would make this sublanguage context-free, a contradiction.

The seed claim is Proposition 34.2. For the algorithm, first test for interior rays. If any exist, test every pair for unequal rational velocities and (65). These arithmetic operations have polynomial bit complexity. Finite points do not affect the language class. $\square$

Corollary 37.2 (Classification for every finitely stabilizing semilinear one-marker seed). *Given a semilinear R , finite stabilization of L_R is decidable. When it occurs, the language-class alternatives of Theorem 37.1 are decidable after effective ray normalization. In particular, every context-free sparse hull in this family is linear context-free. The same conclusions hold from a CFG promised to describe a binary seed containing 0^* and having at most one one per word: its semilinear mask is exactly P_1^L .*

[write, batch 66] For this family of seeds, Theorem 37.1 and Corollary 37.2 answer Part I's Question 11.4 completely, through a restriction on the position support $P_1^L = E_1$ of the kind that question suggests: a finitely stabilizing semilinear one-marker seed has a context-free hull exactly when no two length-compatible interior rays have distinct velocities. Part I's seed S_q is not covered: its mask has components with non-collinear periods, and $K(R) = \infty$ (as $k_R(qm + r) = 2m$), so its non-context-freeness (Theorem 9.2) lies outside the sparse case.

Example 37.3 (Different velocities need not interact). Take rays with bases $(0, 0)$ and $(0, 1)$ and steps $(1, 3)$ and $(3, 1)$, respectively. Their velocities are $1/4$ and $3/4$, but their lengths are $1 \pmod{4}$ and $2 \pmod{4}$. They never coexist. The hull is linear context-free and nonregular, and $K(R) = 1$, so it is already equal to its seed. Dropping the compatibility condition would give a false non-context-freeness test.

Remark 37.4 (Extension to finitely many marker colors). A fixed neutral letter 0 may be supplemented by several marker colors, with each seed word containing at most one nonzero letter. For a finite union of colored arithmetic rays, the available color set at a site is constant on each sufficiently refined length progression. After choosing sites and colors, the same gap-ray proof works with a fixed finite marker word instead of a string of ones. The language-class test depends on the union of nonzero sites, and exact depth depends on

their maximum number, not on the number of available colors. Pumping a nonzero letter again violates the bounded total number of nonzero letters. This observation does not cover seeds with two or more nonzero letters per donor.

38 An exact compatibility-graph formula

The velocity test uses only pairs of rays. Optimal stabilization depth, by contrast, depends on arbitrarily large compatible collections.

Definition 38.1. For the infinite rays of (64), let Γ_R be the simple graph with an edge between distinct vertices v, w precisely when

$$n_v \equiv n_w \pmod{\gcd(q_v, q_w)} \quad \text{and} \quad f_v \neq f_w \text{ as affine functions.} \quad (68)$$

The clique number of the empty graph is zero. Let $E = \{i + j + 1 : (i, j) \in F\}$ be the finite set of exceptional point lengths.

Theorem 38.2 (Clique formula for maximum support). *If $F = \emptyset$, then*

$$K(R) = \omega(\Gamma_R). \quad (69)$$

In general,

$$K(R) = \max \left(\omega(\Gamma_R), \max_{n \in E} k_R(n) \right), \quad (70)$$

where an empty inner maximum is zero. In particular, with no exceptional points,

$$d_{\parallel}(R) = \lceil \log_2 \max(1, \omega(\Gamma_R)) \rceil.$$

Every clique is realized by distinct available sites on an effectively computable infinite arithmetic progression after a finite threshold.

Proof. Fix a length containing only infinite-ray contributions. Choose one contributing ray for each distinct available position. Any two chosen rays have compatible length congruences. They cannot have identical affine position functions because their positions at this length are distinct. Hence the chosen vertices form a clique, giving $k_R(n) \leq \omega(\Gamma_R)$.

Conversely, let S be a clique. Pairwise consistency of the congruences $n \equiv n_v \pmod{q_v}$ implies their joint consistency by the generalized Chinese remainder theorem. Their common solutions form an infinite arithmetic progression. Take solutions above every starting length n_v . For two vertices in S , their affine functions are distinct; they either never meet or meet at one length. Avoid the finitely many crossings. Every remaining common length has $|S|$ distinct available positions. This proves (69), and the construction is effective using CRT and rational comparisons.

Finite points can affect support only at lengths in E . At all other lengths the preceding bound holds, while arbitrarily large lengths still realize a maximum clique and avoid E . Checking the finitely many exceptional lengths gives (70). $\square$

Remark 38.3 (Why a graph of congruences alone is insufficient). Two rays may have different starting points or length periods but the same affine position function. When they coexist they offer one site, not two. This is why (68) excludes identical functions. By contrast, distinct functions that cross at one length still give an edge: the crossing is a finite exception and cannot change the cofinal maximum. Both cases are included in the executable tests.

Corollary 38.4 (Effective maximizing witnesses). *For a sparse ray input, one can compute $K(R)$, both stabilization depths, and a word with $K(R)$ ones in the hull. For every $g < d_{\parallel}(R)$, one can compute a word witnessing $L_R^{[g]} \neq L_R^{[g+1]}$.*

Proof. Compute a maximum clique by finite search and compare exceptional lengths. CRT and crossing avoidance supply a maximizing length for the clique, or an exceptional length realizes the larger value. Select all sites for the first claim. For strictness at generation g , choose any $2^g + 1$ of the maximizing sites. This count is at most 2^{g+1} , so the resulting word belongs to the next generation but not the current one. $\square$

Computing the clique number may require exponential time. The next result shows that this is not merely an artifact of this particular formula.

39 Exact-depth complexity for succinct regular seeds

All coordinates and thresholds in this section are encoded in binary. The input is a list of rays, not an explicit DFA or an expanded grammar. The phrase “regular seed” refers to the represented language, not to polynomial-size availability of a recognizing DFA.

Lemma 39.1 (Every finite graph is an endpoint-ray graph). *Given a finite simple graph G on N vertices, one can construct, in polynomial time and with polynomially many bits, a union R_G of N vertical arithmetic rays such that*

$$\Gamma_{R_G} \cong G, \quad K(R_G) = \omega(G).$$

Both L_{R_G} and Rec_{R_G} are regular.

Proof. Label the vertices $0, \dots, N-1$. Assign a different prime p_{vw} to each nonedge $\{v, w\}$ with $v < w$. For a vertex v , let m_v be the product of the primes assigned to its incident nonedges, with empty product one. Impose residue zero at the smaller endpoint and residue one at the larger endpoint of each nonedge. The ordinary CRT combines these conditions into one residue $c_v \pmod{m_v}$.

For a nonedge $\{v, w\}$, the two vertex moduli share exactly its private prime, and the assigned residues disagree there. Thus the two congruences are inconsistent. For an edge they share no assigned prime, so their moduli are coprime and the congruences are consistent. Choose an integer $B_v \geq v+1$ with $B_v \equiv c_v \pmod{m_v}$, and use the ray

$$(v, B_v - v - 1) + \mathbb{N}(0, m_v). \tag{71}$$

Its marker position is the constant v , distinct at different vertices, and its length congruence is exactly the one constructed. Therefore the compatibility graph is G , and Theorem 38.2 gives the clique identity.

The seed is a finite union of 0^* and languages $0^v 10^{B_v - v - 1} (0^{m_v})^*$, so it is regular. All rays have velocity zero, so the hull is regular by Theorem 37.1. There are at most $N(N-1)/2$ primes. Standard bounds on the r -th prime, for example those of Rosser and Schoenfeld [18], permit their generation by a polynomial-size sieve. Each product m_v has polynomial bit length, and CRT uses polynomial-time integer arithmetic. A valid B_v is obtained by adding the least necessary multiple of m_v , also with polynomially many bits. $\square$

Theorem 39.2 (Support and depth thresholds). *For binary-encoded finite ray unions:*

- (i) *Deciding $K(R) \geq k$, with k part of the input, is NP-complete.*
- (ii) *Deciding whether the parallel process stabilizes by an input integer $g \geq 0$ is coNP-complete.*

Both hardness statements hold with no finite exceptional points and with all rays vertical, so both seed and full closure are regular.

Proof. For (i), a clique of size k is an NP certificate for the cofinal maximum. With finite exceptional points allowed, an exceptional point length together with direct support evaluation is an alternative certificate. The clique formula proves completeness of these certificates. If k exceeds the number of supplied rays and points, the answer is immediately negative. NP-hardness follows from CLIQUE [17] and Lemma 39.1.

For (ii), stabilization by g is equivalent to $K(R) \leq 2^g$. Its failure has the NP certificates just described, of size $2^g + 1$; when this exceeds the number of input components, failure is impossible. One need not construct an exponentially long binary integer for 2^g : first compare g with the logarithmic bound from the input component count. Thus membership in coNP holds also when g itself is binary.

For hardness, start with a nontrivial CLIQUE instance (G, k) with $2 \leq k \leq |V(G)|$. Put

$$d = \lceil \log_2 k - 1 \rceil, \quad \theta = 2^d + 1, \quad u = \theta - k.$$

Here $0 \leq u < k$. Join G with a clique of u new universal vertices, obtaining G' . Then $\omega(G') = \omega(G) + u$, and therefore

$$\omega(G) \geq k \iff \omega(G') \geq 2^d + 1 \iff d_{\parallel}(R_{G'}) > d.$$

This is a polynomial reduction from CLIQUE to failure of stabilization by an input depth. Its complement is coNP-hard, as required. $\square$

Remark 39.3 (Two distinct complexity boundaries). The structural class of a sparse ray hull uses only pair tests and is polynomial-time decidable, while optimal depth detects an arbitrary clique. There is no contradiction: language class and exact resource optimum are different invariants. Nor does the theorem give coNP-hardness for every fixed depth. For a fixed g , searching all subsets of size $2^g + 1$ is polynomial in the number of rays. Finally, the binary periods in (71) may require very large expanded automata. This result does not establish the optimal complexity of the explicit-DFA stabilization problem studied by the preceding ProveIt report.

[write, batch 66] That explicit-automaton problem is Part I's Question 11.2. Part II, written in parallel, settles it for NFA input (PSPACE-complete over two letters, Corollary 20.3) and places DFA input in PSPACE (Theorem 19.2); the exact DFA complexity remains open (Question 24.1). The two complexity results do not conflict: Part II decides whether stabilization happens at all for an explicit automaton, whereas Theorem 39.2 thresholds the exact depth for a succinct list of rays, whose periods can force exponentially large automata. [Added 1 October 2026, batch 73: Part VI settles the exact complexity for DFA input: finite stabilization of an explicit complete DFA is coNP-complete, already over a binary alphabet (Theorem 67.1). It is a coNP-completeness result too, but for another input and another problem: whether stabilization happens at all, for an explicit DFA, by a 3SAT reduction.]

40 Explicit examples and arbitrary finite depths

40.1 A four-word-per-length non-context-free fixed point

Theorem 40.1 (One crossover suffices). *Let*

$$L = 0^* \cup \{0^t 10^{2t+1} : t \geq 0\} \cup \{0^{2t+1} 10^t : t \geq 0\}. \quad (72)$$

Then L is binary and linear context-free, and

$$T(L) = \text{Rec}(L) = L \cup \{0^t 10^t 10^t : t \geq 0\}. \quad (73)$$

Its exact stabilization depth is one. Its full closure is not context-free, but has four words at lengths $2 \pmod{3}$ and one word at every other length. Its generating function is

$$\sum_{n \geq 0} |\text{Rec}(L) \cap \{0, 1\}^n| z^n = \frac{1 + z + 4z^2}{1 - z^3}. \quad (74)$$

Proof. An explicit linear grammar is

$$\begin{aligned} S &\rightarrow Z \mid A \mid B, \\ Z &\rightarrow 0Z \mid \varepsilon, \\ A &\rightarrow 0A00 \mid 10, \\ B &\rightarrow 00B0 \mid 01. \end{aligned}$$

The two marker rays have bases $(0, 1)$, $(1, 0)$ and steps $(1, 2)$, $(2, 1)$. At lengths $n = 3t + 2$ they offer positions t and $2t + 1$; no other length offers a one. Thus the hull consists of the all-zero word, the two single-marker words, and the word placing ones at both positions. The latter is $0^t 10^t 10^t$. Theorem 34.3 gives (73) and depth one; at $t = 0$, the word 11 already witnesses strictness.

The two velocities are $1/3$ and $2/3$ and are length compatible, so Theorem 37.1 proves non-context-freeness. More directly, intersection with $0^* 10^* 10^*$ gives precisely $\{0^t 10^t 10^t : t \geq 0\}$, a gap ray with three positive slopes, which is not context-free by Lemma 36.1. The counts and the geometric-series sum give (74). $\square$

This example separates two phenomena especially clearly. The process has no long transient: it is finished after one generation. Yet it has already crossed the context-free boundary. Conversely its growth is bounded by four words per length, so even extremely small rational growth does not prevent the loss of context-freeness. The preceding ProveIt example with an unbounded transient addresses a different phenomenon; it is not needed in this proof. [write, batch 66] That example is Part I's S_q (Theorems 9.2 and 9.4). The series (74) and (76) are further instances of the commutative series that Part I's Question 11.6 asks to characterize; they have the form of Theorem 8.3 with $p = 3$ and $p = k + 1$.

40.2 A uniform family with prescribed depth

For $k \geq 1$ set

$$R_k = \bigcup_{j=1}^k ((j-1, k-j) + \mathbb{N}(j, k+1-j)), \quad L_k^{\text{fam}} = L_{R_k}. \quad (75)$$

For each j the marker branch is generated by $A_j \rightarrow 0^j A_j 0^{k+1-j} \mid 0^{j-1} 10^{k-j}$; adjoin 0^* .

Theorem 40.2 (All finite depths and exact generation counts). *For $k \geq 1$, the active lengths are exactly $n = (k+1)t + k$, $t \geq 0$. At such a length there are k available sites, at positions $j(t+1) - 1$ for $j = 1, \dots, k$. Consequently*

$$K(R_k) = k, \quad d_{\parallel}(R_k) = \lceil \log_2 k \rceil, \quad d_{\text{frozen}}(R_k) = k - 1.$$

Put $B_{k,g} = \sum_{r=0}^{\min(k, 2^g)} \binom{k}{r}$. Then

$$\sum_{n \geq 0} |L_k^{\text{fam}[g]} \cap \{0, 1\}^n| z^n = \frac{1}{1-z} + (B_{k,g} - 1) \frac{z^k}{1-z^{k+1}}. \quad (76)$$

The full closure replaces $B_{k,g}$ by 2^k . For $k = 1$ the seed is already a linear nonregular fixed point. For every $k \geq 2$, all generations from one onward are non-context-free. Every positive integer d is the exact stabilization depth of such a binary linear seed, for instance with $k = 2^{d-1} + 1$.

Proof. The j -th ray has coordinate sum $k - 1 + (k + 1)t$, so its word length is $(k + 1)t + k$ and its left coordinate is $j - 1 + jt = j(t + 1) - 1$. These positions are distinct. The depth and counts follow from Theorem 34.3 and Corollary 34.4. Summing the count excess over the active arithmetic progression gives (76). The velocities are $j/(k + 1)$, and all rays have the same length progression. For $k = 1$ there is one interior velocity; for $k \geq 2$ there are two distinct co-occurring interior velocities. Theorem 37.1 completes the proof. $\square$

For example, $k = 5$ has seed counts six at each active length, then generation counts 16, 31, and 32. The process stabilizes at depth three, although non-context-freeness appears already at depth one. The closure count does not disclose at which generation the last support combinations became available.

40.3 A regular seed with the same quantitative difficulty

Take the cycle graph C_5 and apply Lemma 39.1. The output rays all have fixed left positions, so the seed and hull are regular. Their maximum support is $\omega(C_5) = 2$, hence depth one. Replacing C_5 by any finite graph encodes its clique number without leaving the regular language class. The supplied example file `endpoint_cycle5.json` records the five explicit rays and the computed result. This is a compact demonstration of the succinctness qualification in Theorem 39.2. [write, batch 66] The example file is shipped as `data/03-closure-classification-endpoint_cycle5.json`, with the result in `data/03-closure-classification-endpoint_cycle5.result.json`.

41 Verification, implementation, and proof status

41.1 What the executable component does

The package contains an exact standard-library Python implementation for finite unions of rays and points. A ray is specified by nonnegative integer pairs `base` and `step`; a zero step denotes one finite point. The analyzer computes the structural language class, compatibility graph, maximum support, both stabilization depths, and a maximum clique. For a non-context-free case it returns a pair of distinct interior velocities whose length congruences are compatible, including their common congruence. The returned indices refer to the list of infinite rays, with finite points omitted.

The implementation uses arbitrary-precision integers, exact rational numbers, gcd, generalized CRT, and an exact branch-and-bound maximum clique search. The class test is polynomial in the ray description; the clique routine has exponential worst-case running time, consistent with the hardness theorem. It is a reference analyzer rather than a claim of optimal practical clique software.

```
python code/crossover_masks.py examples/two_markers.json
python code/crossover_masks.py examples/endpoint_cycle5.json
python code/verify.py
```

[write, batch 66] In this report the files named above are shipped as

- `code/03-closure-classification-crossover_masks.py` (the analyzer),
- `code/03-closure-classification-verify.py` (the verifier),
- `data/03-closure-classification-<name>.json` (the example inputs `examples/<name>.json`), and

- `data/03-closure-classification-<name>.result.json` (their recorded analyses).

Do not run the verifier in place: it imports the analyzer under its delivered module name `crossover_masks`, and it always writes `../data/verification.json` relative to its own directory, which in this report is Part I’s receipt. Section 45 and the report’s README give a copy-based recipe; the analyzer itself imports nothing of the package and can be run in place on a shipped example, its result going to standard output or to an explicit `--output` path outside the report.

41.2 Exact finite audit

The verification script was executed successfully with Python 3.13.5. It recorded 12,867 checks, grouped below. All checks passed. Randomized test generation uses the fixed seeds 20260930 and 61030; every check itself uses exact integer arithmetic and exact finite sets.

Audit group	Checks
Generalized CRT, all moduli through 12 and all residue pairs	6,084
Basic edge cases	5
Uniform-family supports, literal generations, counts, depths, classes	288
Depth-one example, exact set comparisons and counts	93
Random ray unions: clique formula versus complete-period scan	600
Random ray unions: pair criterion versus all length phases	600
Random ray unions: literal crossover generations	1,800
All labeled graphs on 1–5 vertices: graph, clique, regularity checks	3,297
Larger random graph reductions	100
Total	12,867

The complete-period scan explicitly goes past starting thresholds and all affine crossings, then covers a full common length period. Thus it checks more than a heuristically chosen short prefix for the finite ray instances tested. The 1,099 labeled graphs on one through five vertices are all checked, not sampled. Each contributes three tests: graph realization, preservation of the clique number, and regularity of the endpoint-ray instance.

41.3 What has not been implemented or certified

The code does *not* implement a general CFG-to-Presburger compiler, a Presburger quantifier eliminator, arbitrary semilinear normalization, or the DFA synthesis in Theorem 33.1. Those are proved algorithms using classical effective constructions, not features of the ray analyzer. The executable tests do not decide context-freeness of arbitrary grammars; they exercise the proved sparse-ray criterion.

No Lean or Rocq formalization was written or built for this manuscript. The 12,867 passing finite checks are not proofs of universal statements. They are independent exact audits of implementations and finite specializations; the universal conclusions depend on the mathematical proofs above. In particular, testing graph realizations through five vertices does not replace the all-graphs CRT construction.

41.4 A practical formalization route

A staged formalization could begin with finite words and the product hull, then establish the one-count generation theorem and binomial slice counts. The next layer is finite arithmetic: generalized CRT, affine

crossing avoidance, compatibility graphs, and the maximum-support theorem. These parts avoid general context-free grammar theory.

The language-class proof requires a separate finite-gap-ray library and an explicit pumping lemma interface. Its crucial statement is the finite-union obstruction of Lemma 36.1; formalizing only a single obstructing sublanguage would not suffice. The global-regularity algorithm then needs effective Parikh semilinearity, Presburger quantifier elimination, unary periodic language synthesis, and finite-state transductions. The repository’s Presburger development is a natural possible bridge, but no such integration is claimed here.

[write, batch 66] That development is `Logic/PresburgerArithmetic`, described in Section 10: Cooper quantifier elimination in Lean, with the decision procedure `PresburgerArithmetic.Formula.presburgerArithmetic_decidable`. It is not connected to Part III; placement beside it confers no formal status, and no statement of Part III is formalized. No ProveIt Lean or Rocq file defines context-free grammars, Parikh images, semilinear sets or crossover.

42 Further research questions

The following are proposed research directions. They are not all asserted to be globally new open problems; each states a precise next target left unresolved by this manuscript.

Question 42.1 (The nonsparse context-freeness problem). Given a Presburger relation $R \subseteq \mathbb{N}^2$ with unbounded diagonal sections, is context-freeness of Rec_R decidable? The sparse criterion uses a bounded number of markers and finite gap rays, which fail here. The example $R = \mathbb{N}^2$ already shows that unbounded support does not force non-context-freeness. A useful first target is a finite union of rank-two linear sets with a fixed small number of periods.

[write, batch 66] This continues Part I’s Question 11.4 beyond the sparse case answered in Section 37; Part I’s seed S_q (Theorem 9.2) is a non-context-free nonsparse instance.

Question 42.2 (Complexity from an actual grammar). What is the optimal complexity of deciding regularity of $\text{Rec}(L)$ from a CFG for L ? Theorem 33.1 proves decidability but allows large intermediate representations. Determine sharp bounds for linear grammars, unambiguous grammars, and grammars whose position supports admit short quantifier-free formulas. Monadic-decomposition bounds cannot be transferred without accounting for this conversion.

Question 42.3 (Explicit automata versus succinct ray descriptions). What is the optimal finite-stabilization complexity for explicit DFA and NFA seeds? The preceding report supplies decidability machinery, while Theorem 39.2 gives a lower bound for binary-encoded ray descriptions. Closing the gap requires either a reduction with small explicit automata or an improved automata algorithm, not merely expanding the enormous periods of our construction.

[Added 30 September 2026, batch 66: Part II answers the NFA case: finite stabilization of NFA seeds is PSPACE-complete, already over a binary alphabet (Corollary 20.3). For DFA seeds it gives a PSPACE upper bound (Theorem 19.2) and a safe-core criterion (Theorem 21.2), improving Part I’s EXSPACE bound; the exact DFA complexity is open, and is Part II’s Question 24.1. What remains of this question is that DFA case and the gap between explicit automata and the succinct ray inputs of Theorem 39.2.] [Added 1 October 2026, batch 73: Part VI settles the exact complexity for DFA input: finite stabilization of an explicit complete DFA is coNP-complete, already over a binary alphabet (Theorem 67.1). What remains is the gap between explicit automata and succinct ray inputs.]

Question 42.4 (Minimum regular representations and certificates). When a full closure is regular, how large must its minimal DFA be in terms of support row types, column types, and the original grammar? Can a minimal or near-minimal symbolic automaton be synthesized without materializing every period class? In the negative case, can the progression certificate of Theorem 33.2 be shortened or canonically chosen? Recent decomposition-size lower bounds make output representation a central issue [15].

Question 42.5 (Arithmetic restrictions on the compatibility graph). Which restrictions on the ray moduli make exact depth tractable? Promising cases include boundedly many prime factors, a divisibility chain of moduli, a fixed modulus, unary-encoded coordinates, and bounded treewidth of the compatibility graph. The graph realization shows that no unrestricted graph-theoretic simplification is available, but special arithmetic classes may force perfect or otherwise tractable graphs.

Question 42.6 (Approximation and parameterized depth). Can useful approximation or fixed-parameter guarantees be obtained for $K(R)$ or $d_{\parallel}(R)$ under natural arithmetic restrictions? The logarithm compresses the optimum but retains exact threshold hardness. Distinguish approximating $K(R)$ multiplicatively from approximating depth additively, and analyze the role of the largest number of simultaneously compatible periods.

Question 42.7 (More than one marker per donor). For seeds with at most $b \geq 2$ nonzero letters per word, the full hull still depends only on one-coordinate supports, but generation depth no longer does. What finite combinatorial invariant replaces the one-count formula? Even two-marker donors can impose geometric constraints on which groups of sites are inherited by one aligned fragment. An exact classification for a fixed number of arithmetic donor templates would connect the sparse theory here to the general aligned-fragment rank of the preceding report.

[write, batch 66] The preceding report's rank is Part I's ρ_L (Definition 3.1); for one-marker seeds it equals the number of ones (Theorem 34.3 and the remark after it).

Question 42.8 (Synthesis of small obstructions). For sparse masks, a pair of compatible interior velocities certifies non-context-freeness. Can such a pair be transformed automatically into a smallest or simplest explicit non-context-free exact-two-marker sublanguage, with a human-readable pumping certificate? For general Presburger masks, is there a similarly bounded family of geometric obstructions, or can every fixed-size obstruction test fail?

Question 42.9 (Proof-carrying arithmetic language analysis). Can the ray analyzer emit certificates checkable by a small verified kernel: CRT witnesses, maximum-clique upper and lower bounds, finite exception checks, and the exact depth? Maximum-clique lower bounds are short; matching upper bounds may not be. A useful architecture would separate a fast untrusted solver from exact certificate checking, and would expose nonregular and non-context-free certificates independently of the difficult optimum calculation.

43 Conclusion

The coordinate hull retains much more decidable structure than an arbitrary language description. For context-free seeds, its genuine position supports are Presburger relations, and this makes full-closure comparison and regularity effective. The regularity mechanism is the classical finite-row principle, used here with a precise constructive language interface.

For one-marker seeds, the interaction between syntax and generation is especially sharp. Uniformly bounded support is exactly finite stabilization. In that regime, endpoint geometry and at most one co-occurring interior velocity are precisely the regular and linear context-free cases; two distinct co-occurring interior

velocities cause an immediate context-free failure. At the same time, the exact transient length is controlled by a clique number and is computationally hard even inside the regular endpoint case.

These results separate growth, language class, and stabilization depth rather than treating them as proxies for one another. They also provide concrete next steps: larger donor patterns, nonsparse arithmetic masks, explicit-automata complexity, and formally checked arithmetic certificates. [Added 30 September 2026, batch 66: Of these, the explicit-automata complexity has since been settled for NFA input in Part II; see the note after Question 42.3.] [Added 1 October 2026, batch 73: And for DFA input in Part VI (coNP-complete).]

44 Proof audit and boundary cases

[write, batch 66] This section and the next were the manuscript's Appendices A and B.

Empty words and empty support. The hull preserves the seed's empty-word bit. In the one-marker model, 0^* supplies ε and all-zero words at every length. If R is empty, $K(R) = 0$ and depth is zero, not a logarithm of zero.

Finite points and short lengths. A finite point may increase the exact maximum support at its length; formula (70) checks those lengths separately. Finite changes to a language do not affect membership in the regular, context-free, or linear context-free classes, so they do not affect the velocity trichotomy. The class proof nonetheless keeps finite exceptions explicitly in its normal form.

Coincident rays and crossings. Identical affine functions count as one position when simultaneously active, even with different congruence domains. Distinct affine functions may coincide at one exceptional length; they are eventually distinct along a common progression. The graph definition handles both without assuming all positions are distinct at all lengths.

No inference from a bad subset alone. A language can contain a non-context-free sublanguage and still be regular. The proofs use intersection with a regular marker-count slice, then the finite-union gap-ray pumping lemma. This is why additional rays cannot hide the obstruction in the sparse setting.

No inference from mere velocity inequality. Different interior velocities must occur at common lengths. The CRT condition is necessary; Example 37.3 is an explicit counterexample when it is omitted.

No accidental complexity transfer. The NP/coNP statements use binary ray input and a variable threshold. They do not imply the same bounds for expanded grammars, explicit DFAs, or a fixed generation. The general CFG regularity theorem asserts only decidability and effective synthesis.

A note on complete-period testing. A finite prefix alone cannot validate an arithmetic-ray maximum: a compatible collection may first coexist at a very large CRT solution. The verification script derives a common period and a threshold past starts and crossings for each small sampled instance, then scans a complete period. The symbolic clique proof is what handles arbitrarily large inputs without that scan.

45 Package and reproducibility

Path	Contents
article.tex, article.pdf	This standalone source and its compiled article; bibliography is internal.
code/crossover_masks.py	Exact ray analysis, graph realization, and uniform marker family.
code/verify.py	Reproducible exact finite audit.
data/verification.json	Machine-readable result and per-group counts.
data/verification.txt	Captured console output of the run.
examples/	Small ray inputs and their recorded analysis.
notes/provenance.md	Repository pin, literature sources, and bounded novelty statement.
notes/proof-audit.md	Main dependencies and proof obligations.
README.md, Makefile	Build and execution instructions.

The code requires Python 3.10 or later and no third-party packages. The article uses standard pdfL^AT_EX packages; build it with `latexmk -pdf article.tex`. Running the verifier overwrites the two verification receipts only if console output is explicitly redirected to the text receipt; the script itself writes the JSON receipt. Work in a copy to preserve the delivered record. No repository files were edited in preparing this package.

[write, batch 66] *Shipped layout*. The table above describes the delivered package. In this report the files are shipped, byte-identical, as follows; the manuscript and its PDF are printed or replaced by this report, the delivery README by the report’s README, and the checksum list SHA256SUMS was verified and retired.

Shipped path	Delivered as
code/03-closure-classification-crossover_masks.py	code/crossover_masks.py
code/03-closure-classification-verify.py	code/verify.py
code/03-closure-classification-Makefile	Makefile
data/03-closure-classification-verification.json	data/verification.json
data/03-closure-classification-verification.txt	data/verification.txt
data/03-closure-classification-artifact-validation.json	data/artifact-validation.json
data/03-closure-classification-<name>.json,	examples/<name>.json,
...-<name>.result.json	examples/<name>.result.json, for the five names endpoint_cycle5, finite_point_exception, five_markers, incompatible_velocities, two_markers
03-closure-classification-provenance.md	notes/provenance.md
03-closure-classification-proof-audit.md	notes/proof-audit.md

The shipped Makefile still uses the delivered layout (its `check` target runs `code/verify.py`, which in this report is Part I’s verifier) and must not be used from the report directory. The `pdf_sha256` and `tex_sha256` values in the artifact record identify the delivered PDF and source, not this report’s files. To rerun the audit, copy the two programs into a scratch directory under their delivered names and run the verifier there:

```
mkdir -p run/code run/data run/examples
cp code/03-closure-classification-verify.py run/code/verify.py
cp code/03-closure-classification-crossover_masks.py run/code/crossover_masks.py
cd run && python code/verify.py > data/verification.txt
```

When this Part was written the rerun printed PASS with 12,867 checks; its JSON receipt differed from the shipped one only in the recorded Python version (3.14.4 instead of 3.13.5) and, on Windows, in CRLF line endings, and the five example analyses, written with `--output`, reproduced the shipped `.result.json` files byte for byte apart from CRLF line endings.

Part IV. Quadratic finite rank for constrained crossover

[write, batch 71] *Part IV was added on 1 October 2026 in batch 71 of ProveIt’s incoming-report intake, from one manuscript written against Parts I–III as printed above. It answers Part II’s Question 24.2 as to order: the largest finite aligned-fragment rank $F(s)$ of a binary language recognized by an s -state NFA is quadratic, $(s - 3)^2 \leq F(s) \leq (2s + 5)^2$ for $s \geq 4$. Only the lower bound is new; it replaces Part II’s linear lower bound $\lfloor (s - 3)/2 \rfloor$, and the upper bound is Part II’s (Corollary 19.3), imported, not reproved. The leading constant, the exact $F(s)$ and the order for DFAs stay open. Section 46 records the provenance and fixes the notation of the addition. [Added 1 October 2026, batch 73: the order for DFAs is settled in Part VI, the leading constant (one) in Part VII.]*

46 [write, batch 71] Provenance and notation of Part IV

46.1 One manuscript, one addition

Part IV is batch-71 manuscript 01, *Quadratic Finite Rank for Constrained Crossover*, dated 1 October 2026 (8 US-letter pages as delivered). It is an AI-assisted research note: its title page and PDF metadata read “Research note prepared with OpenAI for Vladimir Reshetnikov”, a wording that differs from the “prepared with ChatGPT” of Parts I–III and is quoted here and in the report’s README only. It arrived as the archive quadratic-finite-rank-package.zip (inner directory quadratic-finite-rank) in ProveIt commit aee32ad45 and was placed in commit 427bca743. It was written against ProveIt commit 1bd730779 (in full 1bd730779e8bc44c8976dcd9f38cad86a0d21c26), which postdates the write of Parts II and III (commit 202aafd08); at the pin this report’s `article.tex` had the Git blob ff492c83, unchanged at the placement commit, so “the pinned report” and its “Part II” in the manuscript are this report and its Part II as printed above. The manuscript cites its pin through a link to lines 3092–3110 of that `article.tex`; Question 24.2 occupies lines 3092–3107, and the last two lines of the range already begin Question 24.3.

Its checker, its recorded output and its build script are shipped beside the files of Parts I–III with the prefix 04-quadratic-finite-rank-; its manuscript, delivery README and PDF are not shipped and survive in the arrival commit, and its checksum list SHA256SUMS (6 files, all verified at placement) was retired by the repository’s policy. Section 54 and the report’s README list the shipped names. The package contained no audit or provenance notes; the delivery README’s pin, attribution and non-claims are recorded in this section.

Every result, proof, example, figure, table, limitation and question of the manuscript is printed here, in its order. Its title page and abstract became this section; its Sections 1–8 are Sections 47–54 (manuscript Section n is Section $n + 46$; its main theorem is Theorem 47.1); it has no appendix. Its four unnumbered research questions became Questions 53.1–53.4. Text marked [write, batch 71] was added when the manuscript was written into this report. No theorem was strengthened or weakened.

What Part IV adds, and what it repeats.

- *Answered:* Part II’s Question 24.2 (“Is $F(s)$ linear, quadratic, or of an intermediate order?”), as to order, for NFAs: Theorem 47.1 gives $(s - 3)^2 \leq F(s)$ for $s \geq 4$. The previous lower bound, $\lfloor (s - 3)/2 \rfloor$ for $s \geq 5$, came from Part II’s complete-DFA family (Theorem 22.1); the construction follows the

route Part II's comment on that question anticipated, a long unary transient combined with components that force genuinely separate fragments. The leading constant and the exact $F(s)$ are not determined (Question 53.1). [Added 1 October 2026, batch 73: Part VII proves that the leading constant is one: $(s - 3)^2 \leq F(s) \leq s^2 + 8s$ for $s \geq 4$ (Theorem 76.2).]

- *Imported:* the upper bound $F(s) \leq (2s + 5)^2$ is Part II's (Theorem 14.1(2) and Corollary 19.3); Section 52 only identifies its inputs.
- *Sharpened:* Part II showed the largest finite parallel stabilization index to be of order $\Theta(\log s)$ (comment after Theorem 22.1); Part IV gives $2 \log_2 s + O(1)$, and order $\Theta(s^2)$ for the frozen-source index, where Part II had a linear lower and a quadratic upper bound.
- *New:* the length gate and its last gap (Section 49), the exact rank law of the gate language (Lemma 50.1), the transfer principle (Remark 50.2), the exact depths and slice counts of the family (Corollary 50.3; compare the binomial slice counts of Part III's Corollary 34.4), and the one-initial-state bound $(s - 4)^2 \leq F_{\text{one}}(s)$. Part II's remark before Theorem 14.1 that a single-initial-state convention gives the same results concerns complexity, not F .
- *Not affected:* the lower bound is for NFAs only. For complete DFAs the best lower bound in this report is still Part II's linear one, and the order is open (Question 53.2). Part II's Question 24.1 on the decision complexity for DFA input, Part I's questions and Part III's questions are untouched. [Added 1 October 2026, batch 73: Both DFA questions are answered in Part VI.]
- *Restated:* Proposition 48.1 is Part I's rank calculus (Theorems 3.2, 3.3 and 3.8, Corollary 3.5) in the form of Part II's Proposition 15.2 and Corollary 15.3. It is printed as a marked restatement, with the manuscript's proof as one more route.
- *Credited:* Lemma 49.2 is Sylvester's two-generator Frobenius formula shifted by m ; the manuscript proves it directly without naming it, and a [write, batch 71] note after its proof gives the credit.

46.2 The manuscript's summary and status

Abstract, as delivered (“the ProveIt report” is this report, and the question it answers is Part II's Question 24.2). The optimal finite aligned-fragment rank of an s -state binary NFA has quadratic order. We construct, for every $s \geq 4$, an ε -free NFA with s states, several initial states, universal coordinate hull, and exact maximal rank $(s - 3)^2$. Together with the existing upper bound in the ProveIt report, this gives

$$(s - 3)^2 \leq F(s) \leq (2s + 5)^2, \quad F(s) = \Theta(s^2).$$

The construction couples a classical two-cycle unary length gate with two monochromatic components. At every missing gate length, admissible fragments are exactly the constant intervals. A direct arithmetic proof identifies the last missing length and an alternating word attains it as rank. The least parallel and frozen-source stabilization indices are respectively $\lceil \log_2(s - 3)^2 \rceil$ and $(s - 3)^2 - 1$. This answers Question 24.2 at the pinned repository revision. The upper bound and the fragment-clock framework are credited inputs; no global priority or formal proof-assistant verification is claimed.

Status, as delivered. The manuscript has no separate status box; its abstract ends with the disclaimer above, and its Sections 53 and 54 state its limitations. Its delivery README adds: the finite tests do not prove the all- m theorem, exhaust all bounded-state NFAs, establish the inherited upper bound, or establish global novelty; the note is not a Lean or other proof-assistant formalization; the claimed contribution is the explicit quadratic state-count lower family and the resulting resolution of the pinned repository's asymptotic question; the package makes no claim of first discovery in the worldwide literature and does not settle sharp constants, the exact $F(s)$, or the analogous DFA extremal problem.

46.3 Notation beside Parts I–III

Part IV uses Part I’s operation $\otimes_c$, crossover map $T(L) = L \otimes_c L$, parallel iterates $L^{[k]}$, aligned-fragment rank ρ_L and coordinate hull $\text{Rec}(L)$ (Definitions 2.1, 2.4 and 3.1), Part II’s frozen-source iterates $Y_{k+1} = Y_k \otimes_c L$ (the same sets as Part I’s $L \otimes_c Y_k$), Part II’s supremum rank $R(L)$ (Corollary 15.3) and its extremal function $F(s)$ (Question 24.2), and Part II’s convention: an NFA is ε -free, may have several initial states and need not be complete. Positions are zero based, as in Parts II and III. One symbol was renamed when the manuscript was written in: its gate language L_m is printed as L_m^{gate} , because Part II’s L_m is the one-1 family of Theorem 22.1, whose bound Part IV improves. No normalization changed. The parameter m is *not* renamed, and it is the symbol most easily misread: here $s = m + 2$, whereas Part II’s constants (27) have $m = s + 2$, and the index m of Part II’s family L_m is its rank.

In Part IV	Meaning in Part IV	Not to be confused with
m	size of the length gate; the automaton has $s = m + 2$ states, so $m = s - 2$	Part II’s $m = s + 2$ in (27); the index m of Part II’s family L_m (rank m)
$q_0, \dots, q_{m-1}; z, o;$ u	gate states; the two monochromatic states; the added source of the one-initial-state variant	Part I’s seed parameter q of S_q
$S_m = \langle m, m - 1 \rangle$	numerical semigroup generated by m and $m - 1$	Part I’s seed family S_q ; Part II’s $S_{n,i}$, $S_{r,a}$; Part III’s S_i
H_m	set of gate lengths $\{0\} \cup (m + S_m)$, a set of integers	the coordinate hull, which Part II’s manuscript wrote $\mathcal{H}(L)$ and this report prints $\text{Rec}(L)$
$g_m = (m - 1)^2$	largest integer outside H_m	the Frobenius number of S_m , which is $g_m - m$ (note after Lemma 49.2)
L_m^{gate}	the gate language (77), written L_m in the manuscript	Part II’s one-1 family L_m (Theorem 22.1); Part III’s L_k^{fam} and L_R
$\text{runs}(w)$	number of maximal constant runs of w	Part III’s marker count $k_R(n)$, which counts ones
$R(L), F(s)$	Part II’s supremum rank and its maximum over s -state NFAs	Part I’s $R_L(n)$ (one slice); Part III’s mask $R \subseteq \mathbb{N}^2$
$d_{\parallel}(L), d_{\text{frozen}}(L)$	least stabilization indices of one language	Part I’s depth $d_L(w)$ of one word; Part III’s $d_{\parallel}(R), d_{\text{frozen}}(R)$ are the same notion for L_R
$D_{\parallel}(s), D_{\text{frozen}}(s)$	largest finite stabilization indices over s -state NFAs	Part II’s $D_L(n)$, the largest depth on one slice of one language
$F_{\text{one}}(s)$	$F(s)$ under a single-initial-state convention	$F(s)$ itself, which allows several initial states
T_s, B_s	Part II’s constants, the same values	the operator $T(L)$; Part I’s B_K

The manuscript’s macros for $\mathbb{N}$, ε , Rec , ρ and $\lceil \log_2 \cdot \rceil$ are the report’s; its $\otimes_c$ differs only by an outer `\mathbin`, which changes no spacing; its `runs` was added. Its `lmodern` and `amssymb` packages are not loaded (the report keeps Part I’s `newtx` fonts); the TikZ library `automata` was added for the figure of Section 51.

47 The question and the result

[write, batch 71] “The report” and “the pinned report” of this Part are this report; “Part II” is Part II above. The manuscript’s citation of the pinned report is replaced by internal references.

The report *Aligned Fragments and Constrained Crossover* in the ProveIt repository asks whether the largest finite aligned-fragment rank realizable by an NFA with at most s states is linear, quadratic, or intermediate in

s (Part II, Question 24.2). Its stated bounds leave a linear-to-quadratic gap. The lower bound below closes that gap by converting a long unary transient into many genuinely separate aligned fragments.

We use the report's convention: an NFA is ε -free, can have several initial states, and need not have a transition for every state and letter. Throughout, $\Sigma = \{0, 1\}$ and $\mathbb{N} = \{0, 1, 2, \dots\}$. The state count includes all components of the automaton.

Theorem 47.1 (Quadratic finite rank). *For every $s \geq 4$ there is an s -state NFA recognizing a language L with $\text{Rec}(L) = \Sigma^*$ and*

$$R(L) = (s - 3)^2 < \infty.$$

Consequently, if $F(s)$ is the largest finite $R(L)$ among binary languages recognized by NFAs with at most s states, then

$$(s - 3)^2 \leq F(s) \leq (2s + 5)^2 \quad (s \geq 4).$$

In particular, $F(s) = \Theta(s^2)$.

The construction and its exact rank are proved here. The inequality $F(s) \leq (2s + 5)^2$ is inherited from Part II of the pinned report; Section 52 identifies its quantitative input. The rank and generation formulas recalled next are existing background. They are included with proofs to make the lower-bound argument self-contained.

48 Aligned fragments and generation depth

[write, batch 71] This section restates Part I's Definitions 2.1, 2.4 and 3.1 and Part II's Corollary 15.3 in Part II's zero-based form; it defines nothing new. Proposition 48.1 is stated once, in Part I (Theorems 3.2, 3.3 and 3.8, Corollary 3.5), and restated in Part II (Proposition 15.2, Corollary 15.3); Part III reproves its hull formula (Lemma 31.3). The manuscript's proof below is one more route.

For equal-length words $x, y \in \Sigma^n$ and a boundary $0 \leq i \leq n$, aligned one-point crossover produces $x[0 : i]y[i : n]$ or $y[0 : i]x[i : n]$. We index positions from zero and use half-open intervals. For languages A, B , let $A \otimes_c B$ be all these outputs, with both parents of the same length. Endpoint cuts are permitted. Put

$$L^{[0]} = L, \quad L^{[k+1]} = L^{[k]} \otimes_c L^{[k]}, \quad Y_0 = L, \quad Y_{k+1} = Y_k \otimes_c L.$$

These are, respectively, the parallel and frozen-source iterations.

For a nonempty word w of length n , an interval $[i, j)$ is *admissible* if a seed $v \in L \cap \Sigma^n$ agrees with w on that interval. The rank $\rho_L(w)$ is the minimum number of nonempty admissible intervals partitioning $[0, n)$, or infinity if no partition exists. Set $\rho_L(\varepsilon) = 1$ if $\varepsilon \in L$, and infinity otherwise. Crucially, each seed has the same total length as w , and its fragment occurs at the same coordinates.

The coordinate hull consists of the words of finite rank. Equivalently, for $n \geq 1$,

$$\text{Rec}(L) \cap \Sigma^n = \bigcap_{i=0}^{n-1} \{a \in \Sigma : \exists v \in L \cap \Sigma^n, v_i = a\},$$

and $\text{Rec}(L) \cap \{\varepsilon\} = L \cap \{\varepsilon\}$. Define

$$R(L) = \sup(\{1\} \cup \{\rho_L(w) : w \in \text{Rec}(L)\}).$$

Proposition 48.1 (The established fragment clocks; Parts I and II, restated). *For every language L and every $k \geq 0$,*

$$L^{[k]} = \{w : \rho_L(w) \leq 2^k\}, \quad Y_k = \{w : \rho_L(w) \leq k + 1\}.$$

Both unions over k equal $\text{Rec}(L)$. If $R(L) < \infty$, the least stabilization indices are $\lceil \log_2 R(L) \rceil$ and $R(L) - 1$, respectively.

Another proof (batch-71 manuscript 01). Write $T(L) = L \otimes_c L$. A fragment inherited from one word of $T(L)$ intersects at most two seed-parent pieces. Intersecting a partition with these crossover boundaries gives

$$\rho_L(w) \leq 2\rho_{T(L)}(w).$$

Conversely, pair consecutive pieces in a shortest seed-fragment partition. A crossover of their two seed witnesses at their common boundary supplies the whole paired interval. A final unpaired interval is still admissible because $L \subseteq T(L)$. Hence, with infinity interpreted in the usual way,

$$\rho_{T(L)}(w) = \left\lceil \frac{\rho_L(w)}{2} \right\rceil.$$

Iteration proves the parallel formula, since rank at most one is seed membership.

A frozen-source crossover retains a prefix or suffix of an existing partition and adds at most one seed fragment. It therefore increases the rank bound by at most one. Conversely, for a partition into d seed fragments, choose its seed witnesses $v_1, \dots, v_d$ in order. Start from v_1 and cross with each successive witness at the corresponding partition boundary, preserving the prefix already assembled. This constructs the target in Y_{d-1} . Endpoint cuts allow padding to later generations.

A finite-rank word is eventually generated by either formula. Conversely, singleton intervals witness finite rank exactly for words of the hull. If $R(L) < \infty$, this integer supremum is a maximum unless the hull is empty, which is already covered by the adjoined value one. The threshold formulas give the asserted indices. An equality between adjacent iterates cannot occur earlier: applying the same operator would preserve that equality forever and prevent reaching the full hull. $\square$

49 A length gate with a quadratic last gap

Fix $m \geq 2$. On vertices $q_0, \dots, q_{m-1}$ put the edges

$$q_i \longrightarrow q_{i+1} \quad (0 \leq i < m-1), \quad q_{m-1} \longrightarrow q_0, \quad q_{m-1} \longrightarrow q_1.$$

Label every edge by both 0 and 1, and make q_0 initial and final. The graph is the classical Wielandt digraph: an m -cycle and an $(m-1)$ -cycle sharing a path. Its use here requires no general theorem on primitive matrix exponents. The graph is classical [20, 21]; the relevant return lengths are computed directly.

Let

$$S_m = \langle m, m-1 \rangle = \{am + b(m-1) : a, b \in \mathbb{N}\}, \quad H_m = \{0\} \cup (m + S_m).$$

Lemma 49.1 (Exact return lengths). *The positive lengths of closed walks from q_0 to itself are exactly $m + S_m$. Consequently, the gate accepts precisely the words whose lengths belong to H_m .*

Proof. A positive first-return walk must traverse $q_0, q_1, \dots, q_{m-1}$ and eventually take the edge to q_0 . Before returning, it may traverse the cycle through $q_1, \dots, q_{m-1}$ any number b of times. Its length is therefore

$m + b(m - 1)$. Any positive closed walk is a concatenation of $a \geq 1$ such first-return walks, so its length is $am + b(m - 1)$ with $b \geq 0$. Conversely, every such expression is realized by a first-return walks, placing all b shorter cycles in the first walk. These lengths are exactly $m + S_m$. The empty walk gives length zero. Since each edge has both labels, every binary word of a realizable length labels an accepting walk, and no word of another length does. $\square$

Lemma 49.2 (The last gap). *For every $m \geq 2$, the largest nonnegative integer outside H_m is*

$$g_m = (m - 1)^2.$$

Proof. Suppose $g_m = m + xm + y(m - 1)$ with $x, y \in \mathbb{N}$. Modulo $m - 1$, this implies $x \equiv m - 2 \pmod{m - 1}$, hence $x \geq m - 2$. It follows that

$$m + xm \geq m(m - 1) > (m - 1)^2 = g_m,$$

a contradiction. The same argument is valid when $m = 2$: the modulus is one and the conclusion is $x \geq 0$.

Conversely, let $N \geq g_m + 1$. Divide $N - m$ by $m - 1$:

$$N - m = q(m - 1) + r, \quad 0 \leq r \leq m - 2.$$

Because $N - m \geq (m - 1)(m - 2)$, we have $q \geq m - 2 \geq r$. Therefore

$$N = m + rm + (q - r)(m - 1) \in m + S_m.$$

Thus g_m is absent and every larger integer is present. $\square$

[write, batch 71] *Credit.* Lemma 49.2 is the classical two-generator case of the Frobenius coin problem, shifted by m . For coprime positive integers a, b the largest integer not in $\langle a, b \rangle$ is $ab - a - b$, a result due to Sylvester [22] (see [23] for the history). With $a = m$ and $b = m - 1$ the largest integer outside S_m is $m(m - 1) - m - (m - 1) = m^2 - 3m + 1$ (which is -1 for $m = 2$, where $S_2 = \mathbb{N}$), and a positive n lies outside H_m exactly when $n - m \notin S_m$; hence $g_m = m + m^2 - 3m + 1 = (m - 1)^2$. The manuscript does not name this source; its direct proof above is kept. The finite check of Section 54 and an independent semigroup computation for $2 \leq m \leq 300$ made during the intake of this manuscript agree with the formula.

50 Turning missing lengths into rank

Add two states z, o , each both initial and final, to the gate. State z has only a 0-loop, state o only a 1-loop, and there are no edges between these states and the gate. The initial and final sets are both $\{q_0, z, o\}$. The resulting automaton has exactly $s = m + 2$ states and recognizes

$$L_m^{\text{gate}} = \{w \in \Sigma^* : |w| \in H_m\} \cup 0^* \cup 1^*. \quad (77)$$

This is an ε -free NFA. In particular, its length- n slice is

$$L_m^{\text{gate}} \cap \Sigma^n = \begin{cases} \Sigma^n, & n \in H_m, \\ \{0^n, 1^n\}, & n \notin H_m. \end{cases} \quad (78)$$

The two monochromatic components are essential to the proof: they supply every coordinate at a missing length while forbidding nonconstant fragments there.

Lemma 50.1 (Exact word rank). *The hull of L_m^{gate} is Σ^* . For every nonempty binary word w of length n ,*

$$\rho_{L_m^{\text{gate}}}(w) = \begin{cases} 1, & n \in H_m, \\ \text{runs}(w), & n \notin H_m, \end{cases}$$

where $\text{runs}(w) = 1 + |\{i : 1 \leq i < n, w_{i-1} \neq w_i\}|$ is the number of maximal constant runs. Also $\rho_{L_m^{\text{gate}}}(\varepsilon) = 1$.

Proof. The words 0^n and 1^n supply both letters at every position of every positive length. The empty word is a seed. Thus the hull is universal. If $n \in H_m$, the target itself is a seed, so its rank is one. If $n \notin H_m$, the only seeds of length n are the two constant words. An interval is admissible exactly when it is constant. Every change of letter forces a boundary between fragments, and the maximal constant runs themselves form an admissible partition. This proves both inequalities for the rank formula. $\square$

Proof of Theorem 47.1. By Lemma 49.2, all lengths exceeding g_m belong to H_m and have rank one. For $1 \leq n \leq g_m$, singleton fragments bound every rank by $n \leq g_m$. At the missing length g_m , an alternating word has exactly g_m runs, so Lemma 50.1 gives rank g_m . Thus

$$R(L_m^{\text{gate}}) = g_m = (m-1)^2 = (s-3)^2.$$

This also includes $m = 2$, for which $g_m = 1$ and $L_m^{\text{gate}} = \Sigma^*$. The lower bound for $F(s)$ follows. The inherited upper bound, detailed in Section 52, completes the two-sided inequality and the quadratic-order conclusion. $\square$

Remark 50.2 (A general transfer principle). The same argument works for any cofinite set of lengths H containing zero. For $L = \{w : |w| \in H\} \cup 0^* \cup 1^*$, rank equals one on full slices and run count on missing slices. If g is the largest missing positive length, then $R(L) = \max(1, g)$; if no length is missing, $R(L) = 1$. Thus compact cofinite unary length gates give lower bounds on crossover rank after adding only two states.

Corollary 50.3 (Exact depths and counting). *The least parallel and frozen-source stabilization indices of L_m^{gate} are*

$$d_{\parallel}(L_m^{\text{gate}}) = \lceil \log_2(m-1)^2 \rceil, \quad d_{\text{frozen}}(L_m^{\text{gate}}) = (m-1)^2 - 1.$$

For a missing positive length $n \notin H_m$,

$$\begin{aligned} |(L_m^{\text{gate}})^{[k]} \cap \Sigma^n| &= 2 \sum_{j=0}^{\min(2^k-1, n-1)} \binom{n-1}{j}, \\ |Y_k \cap \Sigma^n| &= 2 \sum_{j=0}^{\min(k, n-1)} \binom{n-1}{j}. \end{aligned}$$

At lengths in H_m , both counts are 2^n at every generation.

Proof. Use Proposition 48.1 and the exact maximal rank. A binary word with $j+1$ runs is determined by its first letter and the j change positions chosen from the $n-1$ interior boundaries. $\square$


[write, batch 71] Here Y_k are the frozen-source iterates of the seed L_m^{gate} .

51 A six state example

Take $m = 4$, so $s = 6$. The gate length set is

$$H_4 = \{0\} \cup (4 + \langle 4, 3 \rangle) = \{0, 4, 7, 8, 10, 11, 12, \dots\}.$$

Its missing lengths are 1, 2, 3, 5, 6, 9; the last is $9 = (4 - 1)^2$. The automaton below makes the state count and both monochromatic components explicit.



At length nine the only seeds are 0^9 and 1^9 . The alternating word 010101010 has rank nine, forcing four parallel generations or eight frozen-source generations. Every word of length at least ten is already a seed. The complete length-nine counts are:

Generation k	0	1	2	3	4	5	6	7	8
Parallel	2	18	186	510	512	512	512	512	512
Frozen-source	2	18	74	186	326	438	494	510	512

The gap between 510 and 512 in the penultimate applicable generation consists of the two fully alternating words.

52 The inherited upper bound and its consequences

Part II of the pinned report proves that an s -state NFA with finite maximal rank satisfies

$$R(L) \leq B_s = (2s + 5)^2. \quad (79)$$

The source uses exact-length forward and backward reachability sets and obtains a common eventual-periodicity threshold

$$T_s = 2(s + 2)^2 + 2(s + 2).$$

Its periodic-interior criterion gives $R(L) \leq 2T_s + 1 = B_s$. The quadratic unary threshold comes from To's corrected arithmetic-progression theorem [8, Section 3], after adding a source and a sink to normalize the unary NFA. The source's relevant statements are its main theorem in Part II, the lemma headed *A common quadratic transient*, the theorem headed *Interior universality criterion*, and the corollary headed *Uniform finite horizon*. This note imports that upper-bound result rather than presenting its proof as a new contribution.

[write, batch 71] These are Theorem 14.1, Lemma 17.1, Theorem 18.3 and Corollary 19.3; the constants are those of (27), where Part II's $m = s + 2$ is not Part IV's gate size.

For fixed s there are only finitely many labelled NFAs up to the stated state bound and convention. Hence the maximum defining $F(s)$ exists. Let $D_{\parallel}(s)$ and $D_{\text{frozen}}(s)$ be the largest finite stabilization indices under the same convention. Proposition 48.1 gives

$$D_{\parallel}(s) = \lceil \log_2 F(s) \rceil, \quad D_{\text{frozen}}(s) = F(s) - 1.$$

Therefore Theorem 47.1 sharpens the quantitative description to

$$D_{\parallel}(s) = 2 \log_2 s + O(1), \quad D_{\text{frozen}}(s) = \Theta(s^2).$$

The leading coefficient two in the logarithmic formula follows from the two-sided quadratic rank bounds, including the ceiling error of less than one.

[write, batch 71] This sharpens Part II's comment after Theorem 22.1, which gave the order $\Theta(\log s)$ for $D_{\parallel}(s)$.

One initial state. If the NFA convention requires a single initial state, add a fresh accepting source u to the constructed automaton. For $a \in \Sigma$, define its a -successors to be the union of the a -successors of the three former initial states; make only u initial and leave all old transitions and final states unchanged. Every nonempty word is accepted exactly when it was before, and accepting u preserves the empty word. This ε -free conversion adds one state. Thus the one-initial-state version satisfies

$$(s - 4)^2 \leq F_{\text{one}}(s) \leq (2s + 5)^2 \quad (s \geq 5),$$

and also has quadratic order. [Added 1 October 2026, batch 73: With Part VII's bound $s^2 + 8s$, which needs no assumption on the initial states (Theorem 76.2), this gives $(s - 4)^2 \leq F_{\text{one}}(s) \leq s^2 + 8s$ and $F_{\text{one}}(s)/s^2 \rightarrow 1$. This is an observation of the write, not of either manuscript.]

53 Attribution and remaining questions

This result resolves an explicit asymptotic question in the inspected repository. It does not identify an exact extremal formula. In particular, the gap between leading constants one and four remains: the result gives $\liminf F(s)/s^2 \geq 1$ and $\limsup F(s)/s^2 \leq 4$, without proving either a limit or a sharp constant. [Added 1 October 2026, batch 73: Part V proves $\limsup F(s)/s^2 \leq 2$ (Corollary 64.2).] [Added 1 October 2026, batch 73: Part VII proves that the limit exists and equals one (Theorem 76.2); the construction here is its lower bound.]

The equal-length fragment framework and its crossover interpretation are not claimed as new. They are developed in the ProveIt report. Closely related represented-interval and one-point-crossover closure machinery appears earlier in Manzoni, Vanneschi, and Mauri [19, Definition 5.2, Theorem 5.3, and Section 5.1]; that paper treats finite populations and logarithmic generation growth. The two-cycle graph is classical [20, 21]. The specific contribution here is the automaton-state lower-bound construction: its two monochromatic components turn the gate's last absent length into exact maximal fragment rank. No exhaustive priority search has been completed.

Several concrete questions remain. [write, batch 71] The manuscript lists them as an unnumbered enumeration; they are printed as numbered questions, in the same order and words; their bracketed titles were added.

Question 53.1 (Leading constant). What is the best leading constant, or the exact function $F(s)$?

[Added 1 October 2026, batch 73: Part V narrows the leading constant to $[1, 2]$: $(s - 3)^2 \leq F(s) \leq 2(s - 1)^2 + 3$ for $s \geq 4$ (Corollary 64.2).] [Added 1 October 2026, batch 73: For complete DFAs, Part VI proves leading constant one: $(s - 7)^2 + 2 \leq F_{\text{DFA}}(s) \leq s^2 + 3$ (Theorem 67.2). Its proof of the upper bound uses determinism and does not apply to NFAs.] [Added 1 October 2026, batch 73: Answered as to the leading constant in Part VII: it is one, $(s - 3)^2 \leq F(s) \leq s^2 + 8s$ for $s \geq 4$ (Theorem 76.2), and over a q -letter alphabet the upper bound is $s^2 + \max\{8, 2q - 4\}s$ (Theorem 81.1). The exact function is known for $s \leq 4$: $F(1), \dots, F(4) = 1, 2, 5, 7$ (Theorem 84.1, computer-assisted for $s = 2, 3, 4$). The exact $F(s)$ for $s \geq 5$ and the optimal additive term remain open.]

Question 53.2 (Order for DFAs). What is the optimal finite-rank order for binary DFAs? The construction here is genuinely nondeterministic: the gate has two outgoing destinations on either letter, and the language uses a union of components. No equally small DFA is established by this argument.

[write, batch 71] For complete binary DFAs with at most s states the bounds in this report are $\lfloor (s - 3)/2 \rfloor$ (Part II's Theorem 22.1, whose language L_m has rank m and a complete DFA with at most $2m + 3$ states, m being that family's index) and $(2s + 5)^2$ (every DFA is an NFA). [Added 1 October 2026, batch 73: Part V lowers that upper bound to $2(s - 1)^2 + 3$, for every NFA and so for every DFA (Corollary 64.2).] [Added 1 October 2026, batch 73: Answered in Part VI: the order is quadratic with leading constant one, $(s - 7)^2 + 2 \leq F_{\text{DFA}}(s) \leq s^2 + 3$ for $s \geq 8$ (Theorems 67.2 and 73.1), by a deterministic variant of the gate with a selector and distinct letters on the two exits of q_{m-1} . Its exact value and lower-order terms are Questions 74.1 and 74.2.]

Question 53.3 (Tightness of the transfer principle). How tightly can a general cofinite unary automaton's last missing length control finite crossover rank through the transfer principle in Remark 50.2?

Question 53.4 (Formalization). Can the lower construction and the report's upper-bound prerequisites be formalized in a proof assistant? This note includes no Lean theorem or machine-checked universal proof.

[write, batch 71] This extends Part II's Question 24.8 to Part IV's construction; no statement of this report is formalized.

The report's separate question about the exact decision complexity for DFA input also remains unaffected. [write, batch 71] That is Part II's Question 24.1. [Added 1 October 2026, batch 73: It is answered in Part VI (Theorem 67.1).]

54 Reproducible finite checks

The accompanying `check.py` uses only the Python standard library. It constructs the transition relation, directly simulates the gate, independently enumerates small seed slices, and compares literal fragment partitions with the exact-length path criterion. The saved `verification.json` records a successful run with:

- 147,677 gate-length comparisons for $2 \leq m \leq 60$ and $0 \leq n \leq 2m^2$;
- 8,184 literal rank computations on every binary word with $2 \leq m \leq 9$ and $0 \leq n \leq 9$;
- 8,222 independent path-based rank computations, including alternating witnesses at the last missing length and the next length for $2 \leq m \leq 20$.

For the last family, the maximum tested witness rank is 361. The path criterion tests a fragment $w[i : j]$ by reaching its start in exactly i steps, following its fixed labels, and reaching a final state in exactly $n - j$ additional steps. It therefore enforces both alignment and total-length matching.

These checks are finite implementation and consistency tests. They do not establish the all- m theorem, exhaust all s -state NFAs, prove the inherited upper bound, determine exact $F(s)$, or verify global novelty. The mathematical proof is in Sections 49–50. Run `python3 check.py` with assertions enabled; do not use Python’s `-O` option. The source file compiles with `pdflatex`; two passes resolve internal references.

[write, batch 71] *Shipped layout*. The paragraphs above describe the delivered package. In this report its files are shipped, byte-identical, as follows; the manuscript and its PDF are printed or replaced by this report, the delivery README by the report’s README, and the checksum list SHA256SUMS was verified and retired.

Shipped path	Delivered as
<code>code/04-quadratic-finite-rank-check.py</code>	<code>check.py</code>
<code>data/04-quadratic-finite-rank-verification.json</code>	<code>verification.json</code>
<code>code/04-quadratic-finite-rank-build.sh</code>	<code>build.sh</code>

The checker prints its record to standard output and writes no file, so it may be run from anywhere, provided its output is not redirected onto the shipped record:

```
py code/04-quadratic-finite-rank-check.py > <scratch>/verification-new.json
```

(`python3` elsewhere; never with `-O`). When this Part was written the rerun, with Python 3.14.4, took about three seconds and printed exactly the bytes of the shipped record; a shell that writes CRLF line endings on Windows makes the comparison hold only after normalizing them. The build script changes to its own directory and runs `pdflatex` twice on `quadratic-finite-rank.tex`, the delivered standalone manuscript, which is printed here as Part IV and is not shipped; it must not be run in `code/`. The “source file” of the last paragraph above is that manuscript; this report is built as its README describes.

Part V. Rational rank slopes for constrained crossover

[write, batch 73] *Part V was added on 1 October 2026 in batch 73 of ProveIt’s incoming-report intake, from one manuscript written against Parts I–III as printed above; it knew Part IV only as a separate manuscript. It answers Part II’s Question 24.3 in all three of its parts: on every live residue r the largest rank $M_L(n)$ of a length- n hull word is $\sigma_r n + O(1)$, with a rational, effectively computable slope σ_r and without refining the period. It also proves that longest-prefix greedy factorization into aligned fragments is optimal, which turns the rank of one word into a linear pass, and lowers Part II’s finite-rank bound $(2s + 5)^2$ to $2(s - 1)^2 + 3$, so that $(s - 3)^2 \leq F(s) \leq 2(s - 1)^2 + 3$ with Part IV’s lower bound. Section 55 records the provenance and fixes the notation of the addition.*

55 [write, batch 73] Provenance and notation of Part V

55.1 One manuscript, one addition

Part V is batch-73 manuscript 33, *Rational Rank Slopes for Constrained Crossover*, dated 1 October 2026 (11 US-letter pages as delivered). It is an AI-assisted research note: its title page and PDF metadata read “Research note prepared with OpenAI for Vladimir Reshetnikov”, quoted here and in the report’s README only. It arrived as the archive `rational-rank-slopes-package.zip` (inner directory `rational-rank-slopes`) in ProveIt commit `f8c3a392a` and was placed in commit `5e4f1263c`. It was written against ProveIt commit `1bd730779` (in full `1bd730779e8bc44c8976dcd9f38cad86a0d21c26`), the pin of Part IV’s manuscript as well; at the pin this report’s `article.tex` had the Git blob `ff492c83`, which holds Parts I–III as printed above, so “the pinned report” and its “Part II” in the manuscript are this report and its Part II. The manuscript cites its pin through a link to lines 3109–3120 of that `article.tex`, which are exactly Question 24.3 and the comment after it. It cites the manuscript of Part IV, *Quadratic Finite Rank for Constrained Crossover*, as a companion note; that note is Part IV here.

Its two checkers, their recorded outputs and its build script are shipped beside the files of Parts I–IV with the prefix `05-rational-slopes-`; its manuscript, delivery README and PDF are not shipped and survive in the arrival commit, and its checksum list SHA256SUMS (8 files, all verified at placement) was retired by the repository’s policy. Section 65 and the report’s README list the shipped names. The package contained no audit or provenance notes; the delivery README’s pin, attribution and non-claims are recorded in this section.

Every result, proof, example, table, limitation and question of the manuscript is printed here, in its order, with one exception: the proof of Lemma 58.1, which repeats Part II’s proof of the same lemma, is replaced by a pointer (the definitions of Section 59, which restate Part II’s, are kept with a pointer). Its title page and abstract became this section; its Sections 1–10 are Sections 56–65 (manuscript Section n is Section $n + 55$; its main theorem is Theorem 62.1); it has no appendix. Its closing sentence of next questions became Questions 65.1–65.4. Text marked [write, batch 73] was added when the manuscript was written into this report. No theorem was strengthened or weakened.

What Part V adds, and what it repeats.

- *Answered:* Part II’s Question 24.3 (“On a bad residue, does $M_L(n)/n$ always have a limit, perhaps after refining the period? Must a resulting slope be rational, and can it be computed effectively?”), in all three

parts, by Theorem 62.1: the limit exists on every live residue of every valid period, with no refinement; it is rational with denominator at most $p(2^s - 1)$; it is the largest reachable cycle mean of an explicit finite graph, hence effective; and it is positive exactly on the residues with an interior obstruction.

- *Sharpened:* Part II’s residue-wise bounded/linear dichotomy (Corollary 18.5) becomes an exact rational slope; Part II’s finite-rank bound $(2s + 5)^2$ (Theorem 14.1) becomes $2(s - 1)^2 + 3$ (Corollary 64.2), by applying Schwarz’s relation-power theorem to Part II’s bound $2t + 1$, so that the leading constant of $F(s)$ lies in $[1, 2]$ rather than Part IV’s $[1, 4]$; Part II’s interval dynamic program for the rank of one word (36) is replaced by a linear pass (Corollary 58.2).
- *New:* optimality of longest-prefix greedy factorization for aligned fragments, for every seed language (Lemma 57.1, a positional form of a classical principle, credited as such); the subset-reset parser and its weighted graph (Proposition 60.1); the reciprocal slopes $1/k$; the two three-state examples showing that neither $t + 1$ nor $\kappa + 1$ replaces $2t + 1$.
- *Restated:* Lemma 58.1 is Part II’s Lemma 16.1 with Corollary 16.2 (pointer; the repeated proof is not reprinted); Section 59 restates Part II’s periodic system; the parity language of Section 63 is Part II’s Example 22.3 with the parities exchanged; the bound $2t + 1$ used in Corollary 64.2 is Part II’s (42).
- *Imported and credited:* Schwarz’s theorem (Theorem 64.1); maximum cycle means and Karp’s algorithm [26]; greedy parsing (Crochemore–Langui–Mignosi, De Luca–Fici); represented-interval machinery (Manzoni–Vanneschi–Mauri); Part IV’s lower family.
- *Not affected:* the leading constant and the exact $F(s)$ (Question 53.1) remain open, now within $[1, 2]$ [Added 1 October 2026, batch 73: the constant is one, by Part VII]; Part II’s Question 24.1 and Part IV’s Question 53.2 on DFAs, Part II’s Questions 24.4 and 24.5 (to which Part V’s denominator bound and Question 65.3 are related), and the questions of Parts I and III are untouched.

55.2 The manuscript’s summary and status

Abstract, as delivered (“the pinned ProveIt report” is this report, its Question 24.3 is Part II’s Question 24.3, the companion note is Part IV, and the slope written λ_r is printed σ_r). For a regular seed language L , let $M_L(n)$ be the maximum aligned-fragment rank among length- n words in its coordinate hull. On every live residue r of any common eventual period of the forward and backward reachability sets, we prove

$$M_L(n) = \sigma_r n + O_L(1), \quad n \equiv r \pmod{p},$$

where $\sigma_r \in [0, 1]$ is rational and effectively computable. The slope is the maximum reachable cycle mean of a finite deterministic subset-reset graph with edge weights zero and one. The key point is that longest-prefix greedy factorization exactly minimizes the number of aligned admissible fragments, so the original maximum-over-words, minimum-over-partitions quantity becomes an ordinary maximum path weight. All terminal subsets are permitted; no period refinement is needed. This resolves the three parts of Question 24.3 in the pinned ProveIt report. The note also gives a linear-pass fixed-word rank algorithm after viable-layer preprocessing, examples with distinct residue slopes and slope $1/k$, and reproducible finite checks. Separately, Schwarz’s classical relation-power bound sharpens the finite-rank estimate to $R(L) \leq 2(s - 1)^2 + 3$. Greedy parsing and cycle-mean methods are classical; no literature-wide priority or proof-assistant verification is claimed.

[write, batch 73] The delivered abstract writes $O(1)$ where Theorem 62.1 writes $O_L(1)$; the constant depends on L in both.

Status, as delivered. The manuscript has no separate status box; its abstract ends with the disclaimer above, and Section 65 states the scope of its finite checks. Its delivery README adds: the finite checks

do not prove the universal slope theorem or Schwarz’s classical theorem; the parser test covers hull targets only and is not an exhaustive check of infinite-rank inputs or of the empty-word branch; greedy parsing is credited to Crochemore, Langiu and Mignosi (2014), with related discussion by De Luca and Fici (2025), cycle means are classical (Karp 1978), and the quadratic lower bound to the companion note; and no global first-discovery claim, formal Lean verification, tight slope-computation complexity, exact eventual affine periodicity or sharp extremal constant is asserted. It also says that the delivered PDF was compiled and every page visually inspected.

55.3 Notation beside Parts I–IV

Part V uses Part II’s notation throughout: the exact-length sets P_i, R_j (32), the viable layers $S_{n,i}$ and letters $A_{n,i}$, a transient and period t, p , the stable sets $P_a^\infty, R_a^\infty, S_{r,a}, C_{r,a}$, live residues, phase-admissible and liftable words, interior obstructions, $M_L(n), R(L)$ and $F(s)$. NFAs are ε -free, may have several initial states and need not be complete; positions are zero based. One symbol was renamed when the manuscript was written in: its slope λ_r , and the maximum cycle mean λ of Lemma 61.1, are printed σ_r and σ , because Part III’s λ_v is the velocity of a ray site, also a rational number in $[0, 1]$ (Section 29.3). No normalization changed.

In Part V	Meaning in Part V	Not to be confused with
σ_r, σ	rank slope on the live residue r ; largest cycle mean (manuscript λ_r, λ)	Part III’s velocity λ_v of a ray site
bad, good	a live residue with, or without, an interior obstruction	the “bad block” of Part I’s Section 5
$\rho(v)$	least number of liftable pieces of a phase-admissible word v	the rank $\rho_L(w) = \rho_L(w)$ of a hull word, which carries the language as a subscript; (86) compares them
$H_{r,a_0}(m), W(m)$	largest $\rho(v)$ over phase words of length m ; largest path weight	Part IV’s set of gate lengths H_m
$G_{r,a_0}, T_{a,c}(Z), N$	the reset graph, its transition, its vertex bound $p(2^s - 1)$	Part III’s graph Γ_R ; the operator $T(L)$ and Part II’s constant T_s
L_n , e.g. L_3, L_4	the length- n slice $L \cap \Sigma^n$	Part II’s one-1 family L_m ; Part IV’s L_m^{gate}
κ, p in Theorem 64.1	index and period of the powers of a relation E	Part II’s period p of (P_i, R_i) , which these imply

The manuscript’s macros for $\mathbb{N}, \varepsilon, \rho$ and `Rec` are the report’s; its macros `\im` and `\runs` are not used in its text and were not carried over (the report’s `\runs` is Part IV’s). Its `lmodern` and `amssymb` packages are not loaded (the report keeps Part I’s `newtx` fonts).

56 The question and the setting

[write, batch 73] “The pinned ProveIt report” and “the report” of this Part are this report as it stood at the pin, that is, Parts I–III; “Question 24.3” is Part II’s Question 24.3, and “the companion note” is the manuscript of Part IV. The manuscript’s citations of both are replaced by internal references.

Question 24.3 of the pinned ProveIt report asks whether the maximum finite rank per length has a limiting ratio on each bad residue, whether a limit must be rational, and whether it is effectively computable (Part II, Question 24.3). The report establishes a bounded-versus-linear dichotomy. It also identifies the central difficulty: rank is a minimum partition cost, while $M_L(n)$ maximizes that cost over targets. A cycle-mean argument requires a justified way to handle this order of quantifiers.

We consider an ε -free NFA

$$\mathcal{A} = (Q, \Sigma, (\Delta_c)_{c \in \Sigma}, I, F), \quad |Q| = s \geq 1,$$

over a finite nonempty alphabet, with an arbitrary set I of initial states. Missing transitions are allowed. For $Z \subseteq Q$, $\Delta_c(Z)$ is the set of c -successors of states of Z ; extend this notation to words by composition. Let L be the accepted language and $L_n = L \cap \Sigma^n$.

For a nonempty target $w \in \Sigma^n$, an interval $[i, j)$ is *admissible* if some $v \in L_n$ satisfies $v[i : j] = w[i : j]$. The aligned-fragment rank $\rho_L(w)$ is the least number of nonempty admissible intervals partitioning $[0, n)$, or infinity if none exists. Set $\rho_L(\varepsilon) = 1$ when $\varepsilon \in L$ and infinity otherwise. The equal total length and the equal coordinates in the witnessing seed are part of this definition.

The coordinate hull is

$$\text{Rec}(L)_n = \prod_{i=0}^{n-1} A_{n,i}, \quad A_{n,i} = \{c \in \Sigma : \exists v \in L_n, v_i = c\} \quad (n > 0),$$

with $\text{Rec}(L)_0 = L_0$. Thus its nonempty words are exactly the words of finite rank. For $n > 0$, define

$$M_L(n) = \max\{\rho_L(w) : w \in \text{Rec}(L)_n\},$$

where the maximum is zero when the hull slice is empty. A nonempty slice gives $1 \leq M_L(n) \leq n$.

The rank definition, coordinate hull, and relation between rank and crossover generation are existing background from the report. [write, batch 73] They are Part I's Definitions 3.1 and 2.4 and Theorem 3.3, in the zero-based form of Part II's Definition 15.1 and (28); $M_L(n)$ is Part II's (Corollary 18.5). Earlier represented-interval machinery for one-point crossover appears in Manzoni, Vanneschi, and Mauri [19]. The contribution pursued here is the explicit rational slope formula for regular seed languages.

57 Why longest aligned fragments are optimal

Lemma 57.1 (Greedy factorization of hereditary intervals). *Fix a word on $[0, n)$. Suppose its family of admissible nonempty intervals contains every singleton and is closed under restriction to nonempty subintervals. Repeatedly taking the longest admissible interval beginning at the current boundary gives a partition with the fewest possible pieces.*

Proof. Let $0 = q_0 < q_1 < \dots < q_k = n$ be any admissible partition, and let g_j be the greedy boundaries, with $g_j = n$ after greedy has finished. Inductively, $g_j \geq q_j$. Suppose $g_{j-1} \geq q_{j-1}$. If $g_{j-1} \geq q_j$, the next greedy boundary is also at least q_j . Otherwise,

$$[g_{j-1}, q_j) \subseteq [q_{j-1}, q_j)$$

is a nonempty subinterval of an admissible interval and is therefore admissible. The longest extension from g_{j-1} reaches at least q_j . This proves the induction. In particular $g_k = n$, so greedy uses at most k pieces. $\square$

Aligned seed fragments satisfy the hereditary hypothesis: restriction of a witnessing seed gives a witness on every smaller interval. Singletons are admissible precisely for targets in the hull. Therefore the lemma applies to every finite-rank target, for arbitrary seed languages, not only regular ones. It also applies to labelled paths through any specified sequence of viable state layers.

This is a positional form of the classical greedy-parsing principle. Crochemore, Langiu, and Mignosi study optimal greedy parsing under suffix-closure conditions for dynamic dictionaries [24]. Their exact dynamic-dictionary convention has additional conditions, so Lemma 57.1 supplies the needed interval-hereditary argument directly. Factorial-language greedy optimality is also stated by De Luca and Fici [25]. We do not claim greedy optimality itself as a new phenomenon.

58 Viable layers and a fixed word algorithm

Let E be the unlabelled transition relation, the union of all Δ_c . Write

$$P_i = E^i(I), \quad R_j = \{q \in Q : \exists f \in F, (q, f) \in E^j\}.$$

For a fixed total length n , define the viable states at boundary i by

$$S_{n,i} = P_i \cap R_{n-i} \quad (0 \leq i \leq n).$$

These are exactly the states visited after i letters in some accepting length- n run.

Lemma 58.1 (The aligned path criterion; Part II, restated). *For $0 \leq i < j \leq n$, a word $x \in \Sigma^{j-i}$ is an admissible fragment on $[i, j]$ if and only if*

$$\Delta_x(P_i) \cap R_{n-j} \neq \emptyset.$$

Equivalently, x labels a path through the successive viable layers $S_{n,i}, \dots, S_{n,j}$. In particular,

$$A_{n,i} = \{c \in \Sigma : \Delta_c(S_{n,i}) \cap S_{n,i+1} \neq \emptyset\}.$$

[write, batch 73] This is Part II's Lemma 16.1 with its Corollary 16.2, and $P_i, R_j, S_{n,i}, A_{n,i}$ are Part II's (32), (33) and (35). The manuscript credits it as background and proves it again by the same argument (extend a path through the viable layers backwards to an initial and forwards to a final state); that proof is not reprinted.

For $w \in \Sigma^n$, $n > 0$, precompute $S_{n,0}, \dots, S_{n,n}$. Start with $Z = S_{n,0}$ and a piece counter $b = 1$. At position i with letter $c = w_i$, compute

$$U = \Delta_c(Z) \cap S_{n,i+1}.$$

If $U \neq \emptyset$, replace Z by U . If $U = \emptyset$, compute

$$U' = \Delta_c(S_{n,i}) \cap S_{n,i+1}.$$

When $U' = \emptyset$, the singleton is unavailable and the rank is infinite. Otherwise increment b and set $Z = U'$. At the end, return b .

The subset Z records all possible ending states of the current greedy fragment. Failed extension closes that fragment immediately before the current letter; the letter is then consumed as the first letter of the next fragment. Because admissibility is hereditary, failure cannot be repaired by extending farther. The parser therefore realizes Lemma 57.1, and its answer is the exact rank. This is a left-to-right parser with a one-letter lookahead at a boundary, after the length-dependent layers have been computed. [write, batch 73] It replaces, for this purpose, the interval dynamic program (36) of Part II, which tests a quadratic number of intervals; the comparison is the write's, not the manuscript's.

Corollary 58.2 (Fixed word rank complexity). *For an s -state NFA represented by Boolean transition matrices and a word of length $n > 0$, exact rank is computable in*

$$O((|\Sigma| + n)s^2)$$

time and $O(ns + s^2)$ auxiliary Boolean storage, excluding the input transition representation. If E is already available, the time is $O(ns^2)$ and the auxiliary storage is $O(ns)$. For $n = 0$, rank is one exactly when $I \cap F \neq \emptyset$, and infinity otherwise.

Proof. Forming E costs $O(|\Sigma|s^2)$. The forward and backward recurrences and their intersections cost $O(ns^2)$. Each letter requires at most two subset images under a transition matrix, each costing $O(s^2)$. Store $O(n)$ subsets of Q and, when needed, the $s \times s$ union matrix E . The empty-word test follows directly from ε -freeness. $\square$

59 The periodic interior

[write, batch 73] This section restates Part II's Section 18: (80) is (37), (81) and (82) are (39) and (40), the nonemptiness and agreement statement is Lemma 18.1 (its two-sentence justification below repeats that proof), and phase-admissible, liftable and interior obstruction are Definition 18.2. Only the words *bad* and *good* for a live residue with and without an obstruction are new; Part II's Question 24.3 uses "bad residue" in this sense.

The recurrence on pairs (P_i, R_i) is deterministic and has at most 4^s possible values. Thus one can compute $t \geq 0$ and $p \geq 1$ such that

$$P_{i+p} = P_i, \quad R_{i+p} = R_i \quad (i \geq t). \quad (80)$$

No minimal period or sharp transient is required for the slope theorem. All phase subscripts below are interpreted modulo p . Set

$$P_a^\infty = P_i \ (i \geq t, i \equiv a \pmod{p}), \quad R_a^\infty = R_i \ (i \geq t, i \equiv a \pmod{p}).$$

A total-length residue r is *live* if $P_r^\infty \cap F \neq \emptyset$. Equivalently, every sufficiently large length in this residue has a nonempty seed slice. A nonlive residue eventually has $M_L(n) = 0$.

Fix a live residue r and define its stable viable states and coordinate alphabets by

$$S_{r,a} = P_a^\infty \cap R_{r-a}^\infty, \quad (81)$$

$$C_{r,a} = \{c \in \Sigma : \Delta_c(S_{r,a}) \cap S_{r,a+1} \neq \emptyset\}. \quad (82)$$

All these sets are nonempty. Indeed, choose an arbitrarily long accepted length in residue r and a sufficiently interior boundary in phase a . An accepting run supplies a state there and a letter at the next position. Moreover, for $n \equiv r \pmod{p}$,

$$S_{n,i} = S_{r,i \bmod p} \quad (t \leq i \leq n - t),$$

and $A_{n,i} = C_{r,i \bmod p}$ whenever both endpoints of that letter lie in this interior.

A word $v = v_0 \cdots v_{m-1}$ is *phase-admissible* from phase a if $v_j \in C_{r,a+j}$ for every j . It is *liftable* if it labels a path through

$$S_{r,a}, S_{r,a+1}, \dots, S_{r,a+m},$$

starting at any state of the first layer. A phase-admissible word that is not liftable is an *interior obstruction*. A live residue is *bad* if such an obstruction exists at some phase, and *good* otherwise. These are the periodic-interior notions used in the source report.

60 The finite weighted reset graph

Fix r and a starting phase a_0 . For $\emptyset \neq Z \subseteq S_{r,a}$ and $c \in C_{r,a}$, put

$$T_{a,c}(Z) = \Delta_c(Z) \cap S_{r,a+1}.$$

Build a letter-labelled weighted graph whose vertices are pairs (a, Z) with $\emptyset \neq Z \subseteq S_{r,a}$. Its initial vertex is (a_0, S_{r,a_0}) . For each permitted letter $c \in C_{r,a}$, the unique outgoing edge with that letter is

$$(a, Z) \longrightarrow \begin{cases} (a+1, T_{a,c}(Z)), & T_{a,c}(Z) \neq \emptyset, \quad \text{weight 0,} \\ (a+1, T_{a,c}(S_{r,a})), & T_{a,c}(Z) = \emptyset, \quad \text{weight 1.} \end{cases} \quad (83)$$

The reset target in the second line is nonempty by the definition of $C_{r,a}$. Retain only vertices reachable from the initial vertex and call the resulting graph G_{r,a_0} . There are at most

$$N = p(2^s - 1)$$

vertices, and every retained vertex has an outgoing edge. All retained vertices are allowed as terminal vertices.

For a nonempty phase-admissible word v , let $\rho(v)$ be the minimum number of consecutive liftable pieces in a partition of v . Liftable intervals are hereditary and all singletons are liftable. Thus the greedy lemma applies.

Proposition 60.1 (Rank equals one plus reset count). *Every phase-admissible word v from a_0 labels a unique path from the initial vertex of G_{r,a_0} , and every such graph path is labelled by a phase-admissible word. If $|v| > 0$, then*

$$\rho(v) = 1 + \text{the weight of its path.}$$

Consequently, if $H_{r,a_0}(m) = \max_{|v|=m} \rho(v)$ and $W(m)$ is the maximum weight of a length- m path from the initial vertex, then

$$H_{r,a_0}(m) = 1 + W(m) \quad (m \geq 1). \quad (84)$$

Proof. The transition rule is deterministic for each permitted letter and permits every phase-admissible letter sequence. Its current subset is precisely the possible endpoints of the current greedy piece. A zero-weight edge extends that piece. A weight-one edge records failure of extension, closes the previous piece, and starts the next piece with the current letter from the entire current viable layer. The first letter never causes a reset because the starting subset is that entire layer. Therefore the number of pieces is one plus the reset count, and greedy optimality makes this the minimum partition cost. Maximizing the equality over words proves (84). $\square$

This identity resolves the order-of-quantifiers issue in the source question. The minimum partition cost is computed deterministically for each target before any maximization occurs. No unjustified interchange of a maximum and a minimum is used.

61 Cycle means with unrestricted endpoints

Lemma 61.1 (Maximum path weight growth). *Let G be a finite directed graph with zero-one edge weights, a fixed initial vertex, and no reachable sinks. Let $W(m)$ be the maximum weight of a length- m path from that vertex, with unrestricted endpoint. If σ is the maximum mean weight of a reachable directed cycle, then*

$$W(m) = \sigma m + O(1).$$

The maximum can be taken over simple cycles. In particular, $\sigma \in [0, 1]$ is rational and effectively computable, and its reduced denominator is at most the number of reachable vertices.

Proof. Let N be the number of reachable vertices. At least one reachable cycle exists, because an infinite walk can be continued without reaching a sink. Delete cycles successively from any length- m path. The remaining simple path has at most $N - 1$ edges. Each deleted cycle has weight at most its length times σ , and the remaining path has weight at most $N - 1$. Hence

$$W(m) \leq \sigma m + N - 1.$$

For the reverse bound, choose a path of length h to a cycle of length ℓ and mean σ . For $m \geq h$, follow that path, take $\lfloor (m - h)/\ell \rfloor$ full turns around the cycle, and take the necessary prefix of the next turn. Nonnegative edge weights give

$$W(m) \geq \sigma(m - h - \ell + 1).$$

The finitely many smaller m can be absorbed into the constant. The final partial turn is allowed because endpoints are unrestricted; there is no divisibility condition on m .

Every closed walk decomposes into simple cycles, and its mean is a length-weighted average of their means. Thus a maximum-mean simple cycle exists. Its mean is an integer divided by its length, at most N , proving rationality and the denominator bound. Finite enumeration of simple cycles proves effectivity; classical cycle-mean algorithms, such as Karp's method with signs reversed, are more efficient on an explicitly given graph [26]. $\square$

The unrestricted endpoint condition is essential to this short proof. Imposing a final state could exclude some lengths and force separate congruence classes. The rank parser imposes no such condition: every subset reached after the last letter is a valid endpoint.

62 Exact rational slopes on the original residues

Theorem 62.1 (Residue-wise rank slope). *Let L be accepted by an s -state ε -free NFA and let t, p satisfy (80). For every live total-length residue r there is an effectively computable rational number $\sigma_r \in [0, 1]$ such that*

$$M_L(n) = \sigma_r n + O_L(1) \quad (n \rightarrow \infty, n \equiv r \pmod{p}). \quad (85)$$

One may take σ_r to be the maximum reachable cycle mean of $G_{r, t \bmod p}$. Its reduced denominator is at most $p(2^s - 1)$. The slope is positive precisely on the bad residues, and zero on the good residues. No refinement of p is needed.

Proof. Fix a sufficiently large $n > 2t$ with $n \equiv r \pmod{p}$, and put $m = n - 2t$ and $a_0 = t \bmod p$. The interval $[t, n - t)$ has the stable viable-layer system beginning at phase a_0 .

Restrict an optimal partition of a full hull word w to this central interval, discard empty intersections, and apply Lemma 58.1. The resulting pieces are liftable. Conversely, a liftable partition of the central word is a partition into admissible aligned seed fragments and can be extended by at most $2t$ singleton pieces. Therefore

$$\rho(w[t : n - t]) \leq \rho_L(w) \leq \rho(w[t : n - t]) + 2t.$$

Every phase-admissible word of length m extends to a full hull word: independently select any available letter at each boundary coordinate. These alphabets are nonempty because the seed slice is nonempty. Maximizing now gives

$$H_{r, a_0}(m) \leq M_L(n) \leq H_{r, a_0}(m) + 2t. \quad (86)$$

Proposition 60.1 and Lemma 61.1 imply $H_{r,a_0}(m) = \sigma_r m + O(1)$. Since $m = n - 2t$, this proves (85), rationality, the denominator bound, and effectivity.

If the residue is good, every phase-admissible word from a_0 is liftable, so no reset occurs and $W(m) = 0$ for every m . Hence $\sigma_r = 0$.

If the residue is bad, choose an obstruction v beginning at a phase a . Append phase-admissible letters until it forms a block B of length divisible by p . From any nonempty subset of $S_{r,a}$, processing the initial v without a reset would give a lift of v from that subset, hence from $S_{r,a}$, which is impossible. Thus each copy of B forces at least one reset. From the initial vertex, a phase-admissible prefix reaches phase a . Repeating B after this prefix gives paths with at least one reset per $|B|$ letters, so the already proved growth formula implies $\sigma_r \geq 1/|B| > 0$. $\square$

For fixed L, t, p , the bounded constants may initially depend on r ; because there are finitely many residues, one can take their maximum. Likewise, different valid transient choices produce the same slope, since (85) identifies it as the limit of the same ratio $M_L(n)/n$ along the residue.

An effective procedure. Compute a repeated pair in the finite process (P_i, R_i) , extract t, p , compute the stable layers and alphabets for each live r , explore the reachable reset graph, and compute its maximum cycle mean. This is a terminating finite procedure. It is not a claim of a tight input-complexity bound: both the period and the subset graph can be large.

63 Examples and interpretation

63.1 Every reciprocal integer occurs

Fix $k \geq 2$ and let L be the binary language avoiding the factor 1^{k+1} . It is regular and its hull is universal: any prescribed singleton can be surrounded by zeroes. More generally, every block of length at most k is an admissible aligned fragment at every position where it fits, since zero padding creates a seed of the required total length. Thus

$$M_L(n) \leq \left\lceil \frac{n}{k} \right\rceil.$$

For the target 1^n , no admissible fragment has length exceeding k , and partitioning into blocks of length at most k attains the bound. Consequently,

$$M_L(n) = \left\lceil \frac{n}{k} \right\rceil \quad (n \geq 1), \quad \sigma = \frac{1}{k}.$$

For $k = 3$, a longest-prefix parser groups 1^{10} as $111 \mid 111 \mid 111 \mid 1$, giving rank four. This supplies genuinely nonintegral slopes and shows that slope denominators are unbounded over all finite automata.

63.2 Different residues can have different slopes

Let

$$L = \{w \in \{0, 1\}^* : |w| \text{ is odd}\} \cup 0^* \cup 1^*.$$

Its hull is universal. Odd slices are already full, so $M_L(n) = 1$ for odd n . At a positive even length, the only seeds are 0^n and 1^n , so admissible fragments are exactly constant intervals. An alternating target needs n

fragments. Hence

$$M_L(n) = \begin{cases} 1, & n \text{ odd}, \\ n, & n > 0 \text{ even}, \end{cases}$$

and the two slopes for the natural period two are zero and one. Thus a single global limit need not exist; the residue-wise formulation is necessary.

[write, batch 73] This is Part II's Example 22.3 with the two parities exchanged (there the even slices are full and the odd ones have linear rank); the manuscript does not cite it. It is printed because it reads the example as a statement about slopes.

63.3 What the bounded error does and does not say

The theorem strengthens a mere two-sided linear estimate to an exact rational leading coefficient. It does not, by itself, prove eventual equality to an affine function, exact periodicity of the bounded error, or a smallest possible error bound. In the reciprocal example, the ceiling already produces a bounded periodic fluctuation. The theorem needs no refinement of the original period to make the ratio converge.

64 A sharper finite rank bound

The same periodic-interior framework gives an improved quantitative bound for the finite-rank case. Define

$$R(L) = \sup(\{1\} \cup \{\rho_L(w) : w \in \text{Rec}(L)\}).$$

The pinned report proves $R(L) \leq (2s + 5)^2$ when $R(L) < \infty$. Its boundary argument actually gives $2t + 1$ for any valid common transient t . [write, batch 73] These are Part II's Theorem 14.1(2) and the bound (42) of Theorem 18.3, which Part II states for any t, p satisfying (37); $R(L)$ is Part II's (Corollary 15.3). Applying a stronger classical relation-power theorem directly on the original state set improves the leading constant from four to two.

Theorem 64.1 (Schwarz's relation-power bound, imported). *For every binary relation E on an s -element set, there exist integers $\kappa, p \geq 1$ such that*

$$\kappa \leq (s - 1)^2 + 1, \quad E^{i+p} = E^i \quad (i \geq \kappa).$$

No irreducibility assumption is required.

This is Theorem 4.3 on printed page 653 of Schwarz [27]. The index is defined on printed page 634 as the least positive exponent whose relation power repeats at a later exponent; multiplying a repeated equality by further powers gives periodicity from that exponent inclusively. The theorem applies to all relations, including reducible ones. The primitive-graph exponent bound alone would not justify this application.

Corollary 64.2 (Leading constant two). *For every s -state ε -free NFA over a finite alphabet,*

$$R(L) < \infty \implies R(L) \leq 2(s - 1)^2 + 3.$$

For the binary extremal function $F(s)$ under the multiple-initial-state convention, the quadratic lower family from the companion note (Part IV, Theorem 47.1) yields

$$(s - 3)^2 \leq F(s) \leq 2(s - 1)^2 + 3 \quad (s \geq 4).$$

Thus the unresolved leading-constant interval narrows from $[1, 4]$ to $[1, 2]$.

Proof. Apply Theorem 64.1 to the unlabelled relation E on the original s states. With $t = (s - 1)^2 + 1$, its power identities imply both $P_{i+p} = P_i$ and $R_{i+p} = R_i$ for $i \geq t$. No source, sink, or initial-state normalization is needed.

If $R(L) < \infty$, Theorem 62.1 shows that no live residue can be bad. For a nonempty hull word of length $n \leq 2t$, singleton fragments give rank at most $2t$. For $n > 2t$, the central interval $[t, n - t]$ is phase-admissible in a good residue, hence liftable and, by Lemma 58.1, one admissible seed fragment. Its two boundaries contribute at most $2t$ singleton pieces. The empty-word convention and the adjoined value one cause no exception. Therefore $R(L) \leq 2t + 1 = 2(s - 1)^2 + 3$.

The companion note (Part IV, Theorem 47.1) constructs an s -state binary NFA with exact finite rank $(s - 3)^2$: an $(s - 2)$ -state two-cycle length gate plus two monochromatic components. Combining its lower bound with the present estimate proves the claim for $F(s)$. $\square$

The statement about constants means

$$1 \leq \liminf_{s \rightarrow \infty} \frac{F(s)}{s^2} \leq \limsup_{s \rightarrow \infty} \frac{F(s)}{s^2} \leq 2.$$

Neither equality of these limits nor a sharp leading constant has been established. [write, batch 73] This improves the bracket $[1, 4]$ of Part IV (Section 53) to $[1, 2]$, for NFAs. The bound $2(s - 1)^2 + 3$ is smaller than Part II's $(2s + 5)^2$ for every $s \geq 1$. [Added 1 October 2026, batch 73: Part VII proves that the limit exists and equals one, with the upper bound $s^2 + 8s$ (Theorem 76.2). That bound is larger than $2(s - 1)^2 + 3$ for $s \leq 11$, so the corollary above remains the better bound for small s .]

64.1 Why a simple one-boundary improvement fails

The coefficient two in the boundary estimate cannot be removed merely by asserting $R(L) \leq t + 1$. Here is an explicit three-state example. Use states 0, 1, 2, initial set $\{0, 1\}$, final set $\{0\}$, and transitions

State	0-successors	1-successors
0	$\{0, 2\}$	$\emptyset$
1	$\{0, 1\}$	$\{0\}$
2	$\{0, 1, 2\}$	$\{2\}$

Here $P_i = R_i = Q$ for all $i \geq 1$, so the minimal common pair transient is $t = 1$ and the period is one. State 2 has both loops, so all stable interior words lift and $R(L) \leq 3$. The exact length-three seed slice is

$$L_3 = \{000, 001, 010, 100\}.$$

The word 111 lies in the hull, but neither of its length-two intervals is admissible. Hence its rank is three and $R(L) = 3 = 2t + 1 > t + 1$.

Even the full relation-power transient κ cannot replace the two-boundary estimate by $\kappa + 1$. For a second example, use initial set $\{2\}$, final set $\{0, 2\}$, and transitions

State	0-successors	1-successors
0	$\{1\}$	$\emptyset$
1	$\{0, 1\}$	$\{1\}$
2	$\{0\}$	$\{2\}$

Its relation powers satisfy $E^3 = E^2 \neq E$, so $\kappa = 2$ and the period is one. Both reachability sets are Q from exponent two onward; state 1 has both loops, giving finite rank at most five. The exact length-four seed slice is

$$L_4 = \{0000, 0010, 1000, 1110, 1111\}.$$

The hull word 0101 has no admissible length-two interval, so its rank is four, greater than $\kappa + 1 = 3$. [write, batch 73] The manuscript claims only $4 \leq R(L) \leq 5$ for this automaton. A literal computation made when the batch was placed (it is not shipped) finds $R(L) = 4$ exactly; nothing in the manuscript depends on the value. These examples do not disprove a state-count leading constant of one; they rule out two direct transient-only shortcuts to it.

65 Reproducibility and remaining questions

The accompanying `verify_greedy.py` constructs accepted seed words literally, collects their actual aligned substrings, and computes the minimum partition count by interval dynamic programming. It compares this independent oracle with the viable-layer subset-reset parser. The recorded run covers every binary two-state NFA with nonempty initial and final sets through target length six, followed by 300 seeded random NFAs on three to five states through target length eight. The seed is 20261001.

The successful run contains 2,604 automata, 14,407 nonempty length slices, and 309,574 rank comparisons. Only target words in nonempty coordinate hulls are compared, so these counts do not include a systematic test of infinite-rank inputs or the empty-word convention. The delivered checker prints a JSON receipt and uses only Python's standard library. Run it without the `-O` flag, since its assertions perform the comparisons.

The second checker, `check_boundaries.py`, computes the relation-power and reachability-pair orbits, literally enumerates both displayed seed slices, and verifies the witnesses in Section 64. Its output is saved as `boundary-verification.json`.

These finite checks corroborate the parser and the faithful handling of alignment and total length. They do not prove the asymptotic theorem, exhaust all NFAs, or verify a formal proof assistant development. The universal claims are supported by the proofs above. The report's periodic-interior framework, classical greedy principles, and cycle-mean theory are credited background; the present note claims a resolution of the pinned repository question, with no assertion of global first discovery.

Natural next questions include the following. [write, batch 73] The manuscript lists them in one sentence; they are printed as numbered questions, in the same order and words; their bracketed titles were added.

Question 65.1 (Complexity of the slope). What is the exact complexity of computing σ_r from an NFA?

Question 65.2 (Denominators). Are there sharper denominator bounds in terms of the input structure?

Question 65.3 (Distinguishing period). What is the smallest common period needed to distinguish slopes?

Question 65.4 (The bounded error). Is the bounded error in (85) eventually periodic?

The explicit reset representation provides a finite object for studying these refinements. [write, batch 73] Question 65.3 bears on Part II's Question 24.5, and the denominator bound $p(2^s - 1)$ of Theorem 62.1 on Part II's Question 24.4; neither of those is answered. No statement of this Part is formalized; Part II's Question 24.8 applies to it as well.

[write, batch 73] *Shipped layout*. The paragraphs above describe the delivered package. In this report its files are shipped, byte-identical, as follows; the manuscript and its PDF are printed or replaced by this report, the delivery README by the report's README, and the checksum list SHA256SUMS was verified and retired.

Shipped path	Delivered as
code/05-rational-slopes-verify_greedy.py	verify_greedy.py
data/05-rational-slopes-verification.json	verification.json
code/05-rational-slopes-check_boundaries.py	check_boundaries.py
data/05-rational-slopes-boundary-verification.json	boundary-verification.json
code/05-rational-slopes-build.sh	build.sh

Both checkers print their record to standard output and write no file, so they may be run from anywhere, provided the output is not redirected onto a shipped record (the report's README gives the commands). When the batch was placed both were rerun on a copy with Python 3.14.4: the boundary checker in about a second and the greedy checker in about seventy seconds, each printing the bytes of its shipped record apart from CRLF line endings. The build script changes to its own directory and runs `pdflatex` twice on `rational-rank-slopes.tex`, the delivered manuscript, which is printed here as Part V and is not shipped; it must not be run in `code/`.

Part VI. Finite crossover stabilization for deterministic automata

[write, batch 73] *Part VI was added on 1 October 2026 in batch 73, from one manuscript written against Parts I–III as printed above; it knew Parts IV and V only as separate manuscripts. It answers Part II’s Question 24.1: for an explicit complete binary DFA, finite stabilization is coNP-complete, under either clock. It also answers Part IV’s Question 53.2: the largest finite rank of a complete binary DFA with s states satisfies $(s - 7)^2 + 2 \leq F_{\text{DFA}}(s) \leq s^2 + 3$ for $s \geq 8$, so $F_{\text{DFA}}(s)/s^2 \rightarrow 1$. Section 66 records the provenance and fixes the notation of the addition.*

66 [write, batch 73] Provenance and notation of Part VI

66.1 One manuscript, one addition

Part VI is batch-73 manuscript 29, *Finite Crossover Stabilization for Deterministic Automata*, dated 1 October 2026 and marked “Revised edition with sharp leading rank constant” (12 US-letter pages as delivered). It is an AI-assisted research note: its title page and PDF metadata read “Research note prepared with OpenAI for Vladimir Reshetnikov”, quoted here and in the report’s README only. It arrived as the archive `dfa-crossover-package.zip` (inner directory `dfa-crossover`) in ProveIt commit `f8c3a392a` and was placed in commit `26abf259b`. It was written against ProveIt commit `1bd730779` (in full `1bd730779e8bc44c8976dcd9f38cad86a0d21c26`), the pin of Parts IV and V as well, where this report’s `article.tex` held Parts I–III as printed above (blob `ff492c83`). The manuscript cites its pin through a link to lines 3079–3090 of that `article.tex`; Question 24.1 occupies lines 3080–3083 there, line 3079 is blank, and lines 3084–3090 are the comment after the question. It cites the manuscripts of Parts IV and V as companion notes.

The earlier edition. The same batch delivered an earlier edition of this manuscript, batch-73 manuscript 32 (archive `dfa-crossover-package (1).zip`, 10 pages, also pinned at `1bd730779`). It lacks the manuscript’s Section 7 (Section 73 here, with Theorem 73.1 and Lemmas 73.2 and 73.3), takes its upper bound $2(s - 1)^2 + 3$ from the manuscript of Part V, and concludes only $F_{\text{DFA}}(s) = \Theta(s^2)$, with “the optimal leading constant between one and two” among its open refinements; its abstract also says that cyclic SCCs “split” the viable-phase graph into stripes, which the revised edition corrects to “organize the relevant cyclic part” of it, since vertices over acyclic components lie in no stripe. Nothing of it is missing from the revised edition, and its nine support files are byte-identical to the revised edition’s. It was superseded, is not printed and nothing of it is shipped; it survives in the arrival commit. The revision note `REVISION.md` of the revised edition records these changes and says that the four finite checkers and their results are unchanged.

Its four checkers, their recorded outputs and its build script are shipped beside the files of Parts I–V with the prefix `06-dfa-crossover-`; its manuscript, delivery README, revision note and PDF are not shipped and survive in the arrival commit, and its checksum list `SHA256SUMS` (13 files, all verified at placement; the earlier edition’s 12 likewise) was retired by the repository’s policy. Section 74 and the report’s README list the shipped names. The package contained no audit or provenance notes.

Every result, proof, table, remark, limitation and question of the manuscript is printed here, in its order. Its title page and abstract became this section; its Sections 1–8 are Sections 67–74 (manuscript Section n is Section $n + 66$); it has no appendix. Its sentence of open refinements became Questions 74.1–74.4. Text

marked [write, batch 73] was added when the manuscript was written into this report. No theorem was strengthened or weakened.

What Part VI adds, and what it repeats.

- *Answered:* Part II’s Question 24.1 (“What is the exact complexity of deciding finite stabilization for complete DFAs over a fixed binary alphabet?”): coNP-complete under polynomial-time many-one reductions, for both clocks; nonstabilization is NP-complete; the coNP upper bound holds over every explicit finite alphabet (Theorem 67.1). With it the DFA parts of Part I’s Question 11.2 and Part III’s Question 42.3 are settled. Part IV’s Question 53.2 (“What is the optimal finite-rank order for binary DFAs?”): quadratic with leading constant one, $(s - 7)^2 + 2 \leq F_{\text{DFA}}(s) \leq s^2 + 3$ for $s \geq 8$ (Theorems 67.2 and 73.1).
- *Sharpened:* for complete DFAs the largest finite parallel stabilization index is $2 \log_2 s + O(1)$ and the frozen-source index is asymptotic to s^2 (comment after the proof of Theorem 67.2); before, the report had for DFAs Part II’s order $\Theta(\log s)$ and a linear lower bound for the frozen index.
- *New:* the SCC-stripe criterion for a nonempty safe core (Lemmas 69.1–69.3, Theorem 69.4); the polynomial certificate for nonstabilization (Proposition 70.1); the monochromatic cylinder wrapper around classical length testers (Lemma 71.1); the deterministic gate family (Proposition 72.2); the early inclusion of actual letters in stable alphabets and the joint entry and exit budget (Lemmas 73.2 and 73.3); the bound $s^2 + 3$ (Theorem 73.1).
- *Restated:* Section 68 restates Part I’s and Part II’s definitions, Part II’s constants, its interior universality criterion and its safe-core characterization, and recalls the latter’s proof; the first two assertions of Lemma 72.1 are Part IV’s Lemmas 49.1 and 49.2 for a differently labelled gate, proved again (the common completion is new).
- *Imported and credited:* the prime-residue assignment encoding and clause testers of Stockmeyer and Meyer [28]; selector-to-cycle DFA constructions of Gawrychowski et al. [29]; Schwarz’s relation-power theorem (Part V’s Theorem 64.1); the two-generator Frobenius arithmetic (Part IV’s credit to Sylvester).
- *Not affected:* the NFA questions. The bound $s^2 + 3$ is proved for complete DFAs only; the leading constant for NFAs (Question 53.1) stays in $[1, 2]$ by Part V [Added 1 October 2026, batch 73: it is one, by Part VII]. Part III’s coNP-completeness (Theorem 39.2) concerns another input and problem (note in Section 71).

66.2 The manuscript’s summary and status

Abstract, as delivered (“the pinned ProveIt report” is this report and its question Part II’s Question 24.1). Finite constrained self-crossover stabilization for complete binary DFA seeds is coNP-complete, under either parallel or frozen-source iteration. This answers the exact DFA-complexity question in the pinned ProveIt report. For the upper bound, cyclic strongly connected components organize the relevant cyclic part of the exponentially expanded viable-phase graph into at most s structural stripes. Nonstabilization has a polynomial certificate consisting of one live residue and one permitted exit from every viable stripe. For hardness, classical Chinese-remainder length testers are embedded in a complete binary DFA with monochromatic branches that convert uncovered lengths into unbounded rank. Separately, we construct complete binary DFAs with $s = m + 6$ states and finite maximal rank between $(m - 1)^2 + 2$ and $(m - 1)^2 + 4$. A joint entry/exit state-budget argument proves $R(L) \leq s^2 + 3$ for every finite-rank complete s -state DFA, and consequently

$$(s - 7)^2 + 2 \leq F_{\text{DFA}}(s) \leq s^2 + 3 \quad (s \geq 8),$$

so $F_{\text{DFA}}(s)/s^2 \rightarrow 1$. The leading constant is sharp. The periodic-interior and safe-core criteria are credited prerequisites; CRT coding and cycle testers are classical. Full proofs, finite-check scope, and attribution are supplied without a global novelty or formal-verification claim.

Status, as delivered. The manuscript has no separate status box; its abstract ends with the disclaimer above, and Section 74 states the scope of its checks and its claimed contribution. Its delivery README adds: the lower witness is not asserted to attain the constructed language’s maximal rank; all state counts include the selector and the rejecting sink; the bound $s^2 + 3$ is not asserted for NFAs; the finite tests are implementation and consistency checks that do not substitute for the proofs of the SCC-stripe theorem, all CRT reduction instances, the complexity result or the asymptotic rank bounds, and the new uniform upper bound is proved analytically, not by an exhaustive finite test; no proof-assistant formalization is supplied; and no exhaustive worldwide priority claim is made. It also says that every page of the delivered PDF was visually inspected and that “the previous delivered edition is preserved separately”, which is manuscript 32 above.

66.3 Notation beside Parts I–V

Part VI uses the report’s operation $\otimes_c$, iterates $L^{[k]}$ and Y_k , rank ρ_L , hull $\text{Rec}(L)$ and supremum rank $R(L)$, and Part II’s P_i , R_j , P_a^∞ , R_a^∞ , live residues, $S_{r,a}$, $C_{r,a}$, phase-admissible and liftable words and safe cores, for a complete DFA $(Q, \Sigma, \delta, q_{\text{in}}, F)$. Two symbols were renamed when the manuscript was written in: its deterministic gate language L_m is printed L_m^{dfa} , because Part II’s L_m is the one-1 family and Part IV’s gate language is L_m^{gate} ; and the initial state of its gate automaton, named r in the manuscript’s Section 6, is printed q_{in} , because r is a residue. No normalization changed. The manuscript itself reuses several letters; the table lists them.

In Part VI	Meaning in Part VI	Not to be confused with
$F_{\text{DFA}}(s)$	largest finite $R(L)$ over complete binary DFAs with at most s states	Part II’s $F(s)$, over NFAs; Part IV’s $F_{\text{one}}(s)$
L_m^{dfa}	the language of the deterministic gate automaton with $m + 6$ states (manuscript L_m)	Part II’s family L_m ; Part IV’s L_m^{gate} (nondeterministic)
t, p in (88)	Part II’s T_s and p_s	in Section 73, a Schwarz transient and an enlarged period
H, d_H, z_H, κ_H	a cyclic SCC, its period, reference vertex and cyclic classes	Part IV’s H_m ; Part V’s κ (a relation-power index)
$W(H, h)$	a stripe, named by (H, h)	Part V’s path weight $W(m)$
C, c, d, κ (Section 73)	the SCC of a safe stripe, its size, period and classes	the stable alphabets $C_{r,a}$; c also denotes a letter
d (Sections 71–72)	number of selector bits; the rejecting sink state	the SCC period d of Section 73
K, D, N, p_i, B	numbers of testers, selector depth, variables; primes; maximal prime	Part I’s K ; Part II’s $K_s, D_L(n), p_s, B_s$
$S, H, g, B = g + 1$ (Section 72)	Part IV’s S_m, H_m, g_m ; a completion length	Part II’s B_s
a, b, z, o (Section 72)	states of the gate automaton	phases a, b elsewhere
h, β, u, v, k, A, B	entry and exit lengths and vertices, internal transient, padded boundaries	

The manuscript’s macros for $\mathbb{N}$, ε , ρ , Rec and $\otimes_c$ are the report’s (its $\otimes_c$ carried an outer `\mathbin`, which changes no spacing); its `\NP` and `\coNP` were added, and `\Ldfa` for the renamed language. Its `lmodern` and `amssymb` packages are not loaded.

67 Results and precise conventions

[write, batch 73] “The pinned report”, “the source report” and “the report” of this Part are this report as it stood at the pin, Parts I–III; “the preceding quadratic NFA note” is the manuscript of Part IV, and “the companion note” of Section 73 the manuscript of Part V. The manuscript’s citations of all three are replaced by internal references.

The pinned report asks for the exact complexity of deciding finite stabilization for complete DFAs over a fixed binary alphabet (Part II, Question 24.1). Its NFA result is PSPACE-complete; it leaves the deterministic case open. The first result below classifies that case. The second resolves the related extremal-rank order for complete binary DFAs, which remained open in the preceding quadratic NFA note (Part IV, Question 53.2).

Theorem 67.1 (Decision complexity). *Given an explicitly represented complete binary DFA, deciding whether its language has finite constrained self-crossover stabilization is coNP-complete under polynomial-time many-one reductions. This holds for both parallel and frozen-source iterations. Nonstabilization is NP-complete. The coNP upper bound extends to every explicitly given finite alphabet.*

Theorem 67.2 (Quadratic finite rank with sharp leading constant). *For every $m \geq 2$, there is a complete binary DFA with one initial state and exactly $m + 6$ states whose language L_m^{dfa} satisfies*

$$(m - 1)^2 + 2 \leq R(L_m^{\text{dfa}}) \leq (m - 1)^2 + 4 < \infty.$$

Consequently, if $F_{\text{DFA}}(s)$ denotes the maximum finite $R(L)$ among complete binary DFAs with at most s states, then

$$(s - 7)^2 + 2 \leq F_{\text{DFA}}(s) \leq s^2 + 3 \quad (s \geq 8).$$

In particular, $F_{\text{DFA}}(s)/s^2 \rightarrow 1$ as $s \rightarrow \infty$.

The lower witness in Theorem 67.2 is not asserted to attain the constructed language’s maximum rank. The stated two-unit interval is intentional. All state counts include the initial selector, accepting monochromatic states, and rejecting sink.

A complete DFA is $\mathcal{A} = (Q, \Sigma, \delta, q_{\text{in}}, F)$, where $|Q| = s$, $\delta : Q \times \Sigma \rightarrow Q$ is total, and there are no ε -transitions. Our lower bounds use $\Sigma = \{0, 1\}$. All complexity claims concern the original explicit DFA, rather than an already expanded periodic graph.

68 Rank and the inherited periodic criteria

For equal-length words $x, y \in \Sigma^n$, aligned one-point crossover at $0 \leq i \leq n$ yields either $x[0 : i]y[i : n]$ or $y[0 : i]x[i : n]$. For languages A, B , let $A \otimes_c B$ contain all such outputs with one parent from each language. Endpoint cuts are allowed. The two iterations are

$$L^{[0]} = L, \quad L^{[k+1]} = L^{[k]} \otimes_c L^{[k]}, \quad Y_0 = L, \quad Y_{k+1} = Y_k \otimes_c L.$$

For $w \in \Sigma^n$, $n > 0$, a nonempty interval $[i, j)$ is admissible if a seed $v \in L \cap \Sigma^n$ agrees with w on that interval. Rank $\rho_L(w)$ is the minimum number of such intervals partitioning the target; it is infinite if no partition exists. The coordinate hull $\text{Rec}(L)$ consists of the finite-rank words, with ε included exactly when it is a seed. Set $\rho_L(\varepsilon) = 1$ if accepted and infinity otherwise, and

$$R(L) = \sup(\{1\} \cup \{\rho_L(w) : w \in \text{Rec}(L)\}).$$

The report's established fragment formulas are

$$L^{[k]} = \{w : \rho_L(w) \leq 2^k\}, \quad Y_k = \{w : \rho_L(w) \leq k + 1\}. \quad (87)$$

Thus either iteration stabilizes finitely exactly when $R(L) < \infty$; the least indices are $\lceil \log_2 R(L) \rceil$ and $R(L) - 1$. These are credited background (Part I's Theorems 3.3 and 3.8 and Corollary 3.5; Part II's Proposition 15.2 and Corollary 15.3; Part IV's Proposition 48.1), not new claims of this note. [write, batch 73] This section restates the report's definitions in the manuscript's words; it defines nothing new. Formula (87) is Part II's (29).

Let E be the unlabelled transition relation and define

$$P_i = E^i(\{q_{\text{in}}\}), \quad R_j = \{q : \exists f \in F, (q, f) \in E^j\}.$$

The source report supplies the common transient and period

$$t = 2(s + 2)^2 + 2(s + 2), \quad p = \text{lcm}(1, 2, \dots, s + 2), \quad (88)$$

with $P_{i+p} = P_i$ and $R_{i+p} = R_i$ for $i \geq t$. It derives this polynomial transient from corrected unary automata results; the source dependencies are recorded there. What matters here is that t is polynomial, $\log p = O(s \log(s + 2))$, and every cyclic SCC period divides p . [write, batch 73] These are Part II's constants T_s and p_s of (27) (Lemma 17.1, from To's corrected unary theorem [8]); the sets $P_i, R_j, S_{r,a}, C_{r,a}$ below are Part II's (32), (39) and (40) for a DFA, and the safe family and safe core are Definition 21.1.

For a phase $a \bmod p$, write $P_a^\infty = P_i$ and $R_a^\infty = R_i$ for any $i \geq t$ in that phase. A total-length residue r is *live* if $P_r^\infty \cap F \neq \emptyset$. For live r , set

$$S_{r,a} = P_a^\infty \cap R_{r-a}^\infty, \quad (89)$$

$$C_{r,a} = \{c \in \Sigma : \exists q \in S_{r,a}, \delta(q, c) \in S_{r,a+1}\}. \quad (90)$$

All these viable state sets and coordinate alphabets are nonempty. The labelled phase graph has vertices (a, q) for $q \in S_{r,a}$; a permitted letter $c \in C_{r,a}$ advances to $(a + 1, \delta(q, c))$ if that vertex is viable. A word is *phase-admissible* when each letter belongs to the corresponding $C_{r,a}$; it is *liftable* if some viable starting state reads it without leaving this graph.

A *safe family* is a family $U_a \subseteq S_{r,a}$ closed under every permitted letter:

$$q \in U_a, c \in C_{r,a} \implies \delta(q, c) \in U_{a+1}.$$

The greatest such family is the *safe core*. The inherited results used for the complexity proof are:

1. Finite stabilization holds exactly when every live residue has no phase-admissible nonliftable word.
2. For a DFA, a live residue has no such obstruction exactly when its safe core is nonempty. A nonempty safe core meets every phase.

These are the report's *Interior universality criterion* and *Safe-core characterization* (Part II's Theorems 18.3 and 21.2). To recall why determinism matters in the second statement: if the safe core is empty, every vertex has a permitted continuation that leaves the viable graph. Append a killing continuation for one currently surviving state. Each other state has at most one image, so the surviving-image count drops. After at most s rounds no starting state survives, producing an obstruction. Conversely, a safe vertex at every phase reads every permitted word. [write, batch 73] This recalls Part II's proof of Theorem 21.2; it is not a second route.

For polynomial-time tests of the stable sets, use the representative

$$i_a = t + ((a - t) \bmod p).$$

Boolean matrix powering computes E^{i_a} in polynomial time in s and $\log i_a$. Every phase test below uses only polynomially many such operations, without expanding all p phases.

69 Compressing safe cores into SCC stripes

A *cyclic SCC* is a strongly connected component containing a directed cycle, including a singleton with a self-loop. For each such component H , let d_H be the gcd of its closed-walk lengths. Fix a reference vertex z_H and a cyclic-class function

$$\kappa_H : H \rightarrow \mathbb{Z}/d_H\mathbb{Z}, \quad \kappa_H(z_H) = 0,$$

so that every internal edge $u \rightarrow v$ satisfies $\kappa_H(v) = \kappa_H(u) + 1$. These data are computable in polynomial time from directed distances and the gcd of the edge discrepancies $\text{dist}(u) + 1 - \text{dist}(v)$. We have $d_H \leq |H|$, hence

$$\sum_{H \text{ cyclic}} d_H \leq s. \quad (91)$$

In particular $d_H \mid p$ for the period in (88). [write, batch 73] The delivered source of the last display lost a \rm at a line break and printed the subscript as “ $H \text{ mcyclic}$ ”; it is printed as intended, the sum over the cyclic SCCs.

For $h \in \mathbb{Z}/d_H\mathbb{Z}$, define a *stripe*

$$W(H, h) = \{(a, q) : q \in H, a - \kappa_H(q) \equiv h \pmod{d_H}\},$$

where a ranges over phases modulo p . Although a stripe can contain exponentially many phase vertices, it has a polynomial-size name (H, h) .

Lemma 69.1 (Uniform viability). *For fixed live r , every stripe is either wholly viable or wholly nonviable. The stripe $W(H, h)$ is viable exactly when $z_H \in S_{r,h}$.*

Proof. Let $d = d_H$. Closed-walk lengths at any $q \in H$ contain every sufficiently large multiple of d . Indeed, a finite collection of such lengths already has gcd d , and its additive semigroup contains all sufficiently large multiples of d ; concatenating the corresponding walks realizes them. The finite gcd fact follows, after dividing by d , by choosing one generator b and a representable value in every residue modulo b , then adding sufficiently many copies of b .

For every attained residue modulo d of a path from q_{in} to q , fix one witnessing path and append the long closed walks at q . Hence forward membership at q is eventually d -periodic. Prepending these closed walks to witnesses from q to a final state proves the analogous fact for backward membership. Stable values defined with the larger period p therefore factor through residues modulo d : evaluate them at arbitrarily large representatives of each p -phase. This does not require the smaller period’s transient to equal the chosen t .

For $q, v \in H$, choose internal paths $q \rightarrow v$ of length ℓ and $v \rightarrow q$ of length k . Then $\ell + k \equiv 0 \pmod{d}$ and $\ell \equiv \kappa_H(v) - \kappa_H(q) \pmod{d}$. Concatenating the paths and using eventual d -periodicity gives

$$q \in P_a^\infty \iff v \in P_{a+\ell}^\infty, \quad q \in R_{r-a}^\infty \iff v \in R_{r-a-\ell}^\infty.$$

Their intersection makes viability constant throughout a stripe. Taking $q = z_H$ and phase h gives the claimed test. $\square$

Lemma 69.2 (Internal edges are permitted). *If (a, q) belongs to a viable stripe in H and $\delta(q, c) = v \in H$, then $v \in S_{r,a+1}$ and $c \in C_{r,a}$.*

Proof. The edge advances the cyclic class by one, so its endpoints belong to the same stripe. Apply Lemma 69.1 and (90). $\square$

Lemma 69.3 (One safe vertex forces its stripe). *If a safe family contains one vertex in a cyclic SCC H , it contains that vertex's entire stripe.*

Proof. Lemma 69.2 and safety force every internal labelled path from that vertex to remain in the family, with the corresponding phase advance. Paths from q to v inside H exist at all sufficiently large lengths congruent to $\kappa_H(v) - \kappa_H(q)$ modulo d_H : take one fixed path and insert sufficiently long closed walks at its start. For any desired compatible target phase b , choose such a length congruent to $b - a$ modulo p . This is possible because $d_H \mid p$. Every vertex of the stripe is therefore reached. $\square$

Theorem 69.4 (Stripe criterion). *For a live residue r , the safe core is nonempty if and only if some stripe $W(H, h)$ is viable and has no permitted exit from its original SCC:*

$$(a, q) \in W(H, h), \quad c \in C_{r,a} \implies \delta(q, c) \in H. \quad (92)$$

Proof. For sufficiency, a viable stripe satisfying (92) is itself safe: every internal edge advances the cyclic class and stays in that stripe.

For necessity, consider the original SCCs represented by vertices of a nonempty safe core. Choose a sink among these represented SCCs in the condensation partial order. This need not be a sink SCC of the whole automaton. Every represented phase vertex has a permitted successor, because $C_{r,a}$ is nonempty; safety puts that successor in the core. The sink choice keeps it in the selected component H . Thus H is cyclic, since a singleton without a self-loop could not have such a successor. Lemma 69.3 forces an entire stripe into the core. A permitted exit from that stripe would, by safety, lead to a strictly lower represented SCC, contradicting the sink choice. $\square$

The qualification concerning represented SCCs is essential. For example, an accepting 0-loop in a DFA for 0^* can be safe although its forbidden 1-edge leads to a rejecting bottom component. Safety concerns permitted letters, not all letters at every phase.

70 A polynomial certificate for nonstabilization

Proposition 70.1. *Finite stabilization for complete DFAs over an explicitly given finite alphabet belongs to coNP.*

Proof. A certificate for nonstabilization consists of a live residue $r < p$ and, for every viable stripe $W(H, h)$, a triple (q, a, c) satisfying

$$\begin{aligned} q &\in H, & a < p, & & a - \kappa_H(q) \equiv h \pmod{d_H}, \\ c &\in C_{r,a}, & \delta(q, c) &\notin H. \end{aligned}$$

The verifier computes the cyclic SCCs, their periods and cyclic classes, and all stripe names. It checks that r is live. For each stripe it deterministically tests viability using $z_H \in S_{r,h}$. A nonviable stripe requires no exit witness; for each viable one, it verifies the displayed conditions.

There are at most s stripes by (91). Every guessed phase uses $O(\log p)$ bits, polynomial in the input. Stable membership and permitted-letter predicates are computable by polynomially many Boolean matrix powers at polynomial-bit exponents. All other checks are elementary graph and modular arithmetic. Therefore certificate size and verification time are polynomial.

The stripe criterion says such exit witnesses exist for exactly those live residues with empty safe core. The inherited criteria in Section 68 say that at least one such residue exists exactly when stabilization fails. Hence nonstabilization is in NP, and stabilization is in coNP. $\square$

This certificate avoids expanding the full phase graph. An exit may lead to another viable or even safe region; that does not invalidate the criterion. If any safe region existed, a lowest represented SCC would supply some stripe with no permitted exit, and the certificate would fail for that stripe.

71 Binary DFA hardness via length testers

The arithmetic coding in this reduction is classical. Prime-residue assignment encodings and invalid-assignment and clause-falsification testers appear explicitly in Stockmeyer and Meyer’s proof of Theorem 6.1 [28]. Closely related DFA length constructions, including binary selectors feeding unary cycles, appear in Gawrychowski et al. [29, Theorems 5–6]. We credit those ingredients. The crossover-specific step is the monochromatic/cylinder embedding below, which turns uncovered lengths into high aligned rank. [write, batch 73] Part III’s Theorem 39.2 is also a coNP-completeness result, but for thresholding the exact depth of a succinct list of binary-encoded rays, by a CLIQUE reduction; Part VI decides whether stabilization happens at all for an explicit DFA. The two concern different inputs and problems, and neither implies the other. The manuscript does not cite Part III.

71.1 A polynomial family of unary testers

Start with a 3CNF formula φ in variables $x_1, \dots, x_N$. Remove tautological clauses and repeated literals inside a clause. Add an unused variable when necessary so that $N \geq 1$. Choose distinct primes $p_1, \dots, p_N$, each polynomially bounded in N . The first N primes suffice: their largest value is $O(N \log(N + 1))$, and a sieve constructs them in polynomial time.

A length n encodes a Boolean assignment when $n \bmod p_i \in \{0, 1\}$ for every i , with residue one meaning true. Use the following cycle testers:

- For variable i , a cycle of modulus p_i marks the residues outside $\{0, 1\}$.
- For each clause, a cycle whose modulus is the product of its at most three distinct variable primes marks the residue that makes every literal of the clause false. An empty clause uses modulus one and marks every length.

The Chinese remainder theorem determines the clause residue. Call a tester *active* at n if it marks that total length. If φ is unsatisfiable, every length activates some tester: either its residues do not encode an assignment or they encode an assignment falsifying a clause. If φ is satisfiable, its satisfying assignment gives an infinite arithmetic progression of lengths activating none, modulo $\prod_i p_i$.

Let $K = N + \#\{\text{clauses}\} \geq 1$ and $B = \max_i p_i$. The cycles use at most KB^3 states. Their active-residue sets can be computed by enumerating their polynomially many residues. This construction is polynomial in the formula size. A variable tester with modulus two simply has no active residue, which causes no problem.

71.2 A deterministic binary selector with fragile constant branches

Build a complete binary DFA whose first two letters select one of four cases:

Prefix	Destination
00	Accepting state with a 0-loop; letter 1 rejects forever
11	Accepting state with a 1-loop; letter 0 rejects forever
01	Fixed-depth binary selector for the K testers
10	Rejecting sink

The selector uses $d = \lceil \log_2 K \rceil$ additional bits, with unused codes going to the sink. All tester entries are at the same total depth

$$D = 2 + d.$$

For $K = 1$, no extra selector bits or decision nodes are needed. Every tester is a directed cycle on which both input letters advance one step. If its modulus is h and its desired set of active total-length residues is $A \subseteq \mathbb{Z}/h\mathbb{Z}$, declare cycle position j final exactly when

$$D + j \pmod{h} \in A.$$

This explicit offset ensures that activity depends on total word length, including the selector. Prefix and selector states are nonfinal, and all unspecified transitions go to a rejecting sink with both loops. The selector has $O(K)$ states, so the whole DFA remains polynomial in size.

The only seeds accepted outside tester branches are 0^n and 1^n for $n \geq 2$. If a tester is active at $n \geq D$, its fixed selector prefix can be followed by *every* binary suffix of length $n - D$ to obtain a seed. Thus it contributes a full suffix cylinder.

Lemma 71.1 (Correctness of the crossover embedding). *The constructed DFA stabilizes finitely if and only if φ is unsatisfiable.*

Proof. If φ is satisfiable, there are arbitrarily large $n \geq D$ with no active tester. At each such length the entire seed slice is $\{0^n, 1^n\}$. Both symbols occur at every coordinate, so the hull slice is all binary words, but admissible fragments are exactly constant intervals. An alternating target has rank n . Hence $R(L) = \infty$ and no finite stabilization occurs.

If φ is unsatisfiable, some tester is active at every length. For $n > D$ and any hull target w , its first D letters can be supplied as singleton fragments using 0^n and 1^n . The remaining suffix is one admissible aligned fragment supplied by an active tester cylinder. Thus $\rho_L(w) \leq D + 1$. At $n = D$, omit the empty suffix, obtaining rank at most D . At smaller positive lengths, the singleton bound on existing hull slices gives rank at most $n < D$. The length-zero convention is harmless. Therefore $R(L) \leq D + 1 < \infty$. $\square$

Proof of Theorem 67.1. The map from φ to its complete binary DFA is polynomial and sends unsatisfiable 3CNF formulas exactly to finitely stabilizing languages. This gives coNP-hardness. Proposition 70.1 gives membership. Taking complements yields NP-completeness of nonstabilization. Formula (87) makes the two iteration clocks equivalent for this decision problem. $\square$

72 A deterministic family with quadratic finite rank

The rank construction is independent of the complexity reduction. It reuses the classical two-cycle graph underlying the earlier NFA family, while assigning different letters to its two branching edges and using a deterministic selector. [write, batch 73] The earlier NFA family is Part IV's gate language L_m^{gate} (Section 49). Below, S , H and g are Part IV's S_m , H_m and g_m with the index dropped, and the first two assertions of Lemma 72.1 are Part IV's Lemmas 49.1 and 49.2 (Sylvester's two-generator Frobenius formula shifted by m ; see the note after Lemma 49.2), proved again directly; the common completion length is new.

72.1 The complete automaton

Fix $m \geq 2$ and use gate states $q_0, \dots, q_{m-1}$ with

$$\delta(q_i, 0) = \delta(q_i, 1) = q_{i+1} \quad (0 \leq i < m-1),$$

$$\delta(q_{m-1}, 0) = q_0, \quad \delta(q_{m-1}, 1) = q_1.$$

Only q_0 is final inside the gate. Add exactly six states: an initial state q_{in} , prefix states a, b , accepting monochromatic states z, o , and a rejecting sink d . Their transitions are

State	0-successor	1-successor
q_{in}	a	b
a	z	q_0
b	d	o
z	z	d
o	d	o
d	d	d

The final set is $\{q_0, z, o\}$ and q_{in} is the sole initial state. [write, batch 73] The manuscript names the initial state r ; it is printed q_{in} , the manuscript's own name for an initial state in Section 67, because r denotes a residue elsewhere. This is a complete binary DFA with exactly $m + 6$ states. Prefix 01 enters the gate; outside that prefix, the seeds are exactly $0^n, 1^n$ for $n \geq 2$. In particular the hull consists of every binary word of length at least two; its length-zero and length-one slices are empty.

72.2 The quadratic missing return length

Put

$$S = \langle m, m-1 \rangle = \{am + b(m-1) : a, b \in \mathbb{N}\}, \quad H = \{0\} \cup (m + S), \quad g = (m-1)^2.$$

Lemma 72.1 (Gate lengths and common completion). *The lengths of paths from q_0 to itself are exactly H , whose largest missing nonnegative integer is g . From every gate state there is a path to q_0 of length exactly $g + 1$.*

Proof. The unlabelled gate graph has an m -cycle and an $(m-1)$ -cycle sharing a path. A positive first return to q_0 has length $m + b(m-1)$, with $b \geq 0$. Concatenating first returns gives exactly $am + b(m-1)$ with $a \geq 1$, hence the set $m + S$. Assigning different labels to the two exits from q_{m-1} changes which words label these paths, but not which path lengths exist.

To see that $g \notin H$, an expression $g = m + xm + y(m-1)$ would imply $x \equiv m-2 \pmod{m-1}$ and therefore $m + xm \geq m(m-1) > g$. This includes $m = 2$, where the modulus is one. If $N \geq g + 1$, divide $N - m = q(m-1) + r$, with $0 \leq r \leq m-2$. Then $q \geq m-2 \geq r$, so

$$N = m + rm + (q-r)(m-1) \in H.$$

Thus g is the last missing return length.

The semigroup S contains every integer at least $(m-1)(m-2)$: write such an integer as $q(m-1) + r$, $0 \leq r \leq m-2$, and use $rm + (q-r)(m-1)$. For $i \geq 1$, the path lengths from q_i to q_0 are exactly $(m-i) + S$, by the same first-return decomposition. With $B = g + 1$,

$$B - (m-i) = (m-1)(m-2) + i \in S.$$

For $i = 0$, $B \in H$ was already shown. Hence every gate state reaches q_0 in exactly B steps. $\square$

Proposition 72.2 (Witness and direct finite bound). *The constructed language satisfies $g+2 \leq R(L_m^{\text{dfa}}) \leq g+4$.*

Proof. At total length $n = g + 2$, a word beginning with 01 would require a gate return of length g , which is impossible. Hence the full seed slice is exactly $\{0^n, 1^n\}$. An alternating target lies in the hull and every admissible fragment is constant, giving rank exactly $n = g + 2$.

For the upper bound, put $B = g + 1$. Let w have length $n > B + 2$. Its central interval $w[2 : n - B]$ is nonempty. Read this central word from q_0 ; completeness keeps its unique run inside the gate. Lemma 72.1 supplies a completion of length B from its last state to q_0 . Prefixing 01 produces an accepted length- n seed agreeing with the whole central interval at its proper coordinates. The first two target letters and last B target letters are supplied as singleton fragments by the monochromatic seeds. Consequently

$$\rho_{L_m^{\text{dfa}}}(w) \leq 2 + 1 + B = g + 4.$$

For $2 \leq n \leq B + 2$, singletons give rank at most $n \leq g + 3$. The two shorter hull slices are empty and the baseline value one is below $g + 4$. This proves the uniform finite bound without using a safe-core or complexity theorem. $\square$

Proof of Theorem 67.2. Proposition 72.2 gives the stated family and its finiteness. Set $m = s - 6$ for $s \geq 8$ to obtain $F_{\text{DFA}}(s) \geq (s - 7)^2 + 2$. The complete-DFA upper bound $R(L) \leq s^2 + 3$ is proved in Theorem 73.1 below. Together the inequalities give the asserted bound and, after division by s^2 , the limit one. $\square$

The largest finite parallel stabilization index for complete binary DFAs is therefore $2 \log_2 s + O(1)$, and the largest finite frozen-source index is asymptotic to s^2 , by (87). These are extremal asymptotics; the exact rank or least stabilization index of every individual automaton is not determined here.

73 A sharp leading constant for every complete DFA

The earlier general NFA bound $2(s - 1)^2 + 3$ in the companion note (Part V, Corollary 64.2) charges both boundaries by a global transient. For a DFA, a safe SCC permits a better allocation: the states used before entering that component and after leaving it are disjoint. A separate pumping observation makes the actual target letters compatible with stable alphabets after only s positions at either boundary.

Theorem 73.1 (Uniform deterministic rank bound). *Let L be recognized by a complete s -state DFA over a finite nonempty alphabet. If $R(L) < \infty$, then*

$$R(L) \leq s^2 + 3.$$

The proof uses determinism and does not establish this bound for general NFAs.

[Added 1 October 2026, batch 73: For NFAs Part VII proves $R(L) \leq s^2 + 8s$ over two letters and $R(L) \leq s^2 + \max\{8, 2q - 4\}s$ over q letters (Theorems 76.2 and 81.1), which also gives $F_{\text{DFA}}(s) \leq s^2 + 8s$ for binary DFAs. The point of the theorem here is the additive constant 3 for complete DFAs over every finite alphabet, together with the deterministic lower family.]

For this proof choose a global transient $t \leq (s - 1)^2 + 1$ and a period for E from Schwarz's theorem [27, Theorem 4.3]. Enlarge the period p to a multiple of every integer $1, \dots, s$; this does not change t . The resulting period need not be the particular p used for the complexity certificate. Stable sets, live residues, and stripes are defined as before. The source's interior criterion and the stripe results above apply to any such common transient and period.

Lemma 73.2 (Actual alphabets enter stable alphabets early). *If $i \geq s$, then $P_i \subseteq P_{i \bmod p}^\infty$. If $j \geq s$, then $R_j \subseteq R_{j \bmod p}^\infty$. Consequently, for an actual hull word of total length n , its letter at position i belongs to $C_{n \bmod p, i \bmod p}$ whenever*

$$i \geq s \quad \text{and} \quad n - i - 1 \geq s.$$

The claims are inclusions, not equalities.

Proof. A walk of at least s edges visits a vertex twice and contains a simple directed cycle of length $\ell \leq s$. Since $\ell \mid p$, inserting p/ℓ extra copies of this cycle adds p edges without changing endpoints. Repeating reaches arbitrarily large lengths in the same phase, proving the forward inclusion. Apply the same argument in the reversed graph for the backward inclusion.

For the letter statement, take its accepting seed witness with transition $q \xrightarrow{c} q'$ at position i . Pump the witness prefix and suffix separately. The prefix proves stable forward membership of q , and appending c proves stable forward membership of q' . The suffix proves stable backward membership of q' , and prepending c proves stable backward membership of q . Thus both endpoints lie in the correct stable viable layers and c is permitted there. $\square$

[Added 1 October 2026, batch 73: The proof uses no determinism. Part VII states the lemma for NFAs, with a liveness clause, as Lemma 79.1 and proves it by this argument, which is printed only here.]

Fix a live residue r . Finite rank and the inherited safe-core criterion give a nonempty safe core. Theorem 69.4 supplies an entire viable stripe over an original cyclic SCC C , closed under every stable-permitted letter at its phase. Write $c = |C|$, $d = d_C$, and κ for its cyclic-class function. Full-stripe membership depends only on the class/time offset modulo d , even though the phase graph is indexed modulo the possibly much larger p .

Lemma 73.3 (A joint entry and exit budget). *For that stripe there exist an entry vertex $u \in C$, an exit vertex $v \in C$, an actual initial-to- u path of length h , and an actual v -to-final path of length β , such that*

$$h + \beta \leq (s - c)d. \tag{93}$$

The pair $(h \bmod p, u)$ belongs to the stripe, and for every $n \equiv r \pmod{p}$,

$$n - h - \beta \equiv \kappa(v) - \kappa(u) \pmod{d}. \tag{94}$$

Proof. Stable viability supplies an actual accepting walk through a stripe vertex at its prescribed phase, of total length $N \equiv r \pmod{p}$: choose a sufficiently long prefix and suffix from the stable reachability sets. Let u be its first vertex in C and v its last. Between them the walk cannot leave C and return, since any such outside vertex would be mutually reachable with C and hence belong to it. Let the original entry time and exit length be h_0 and β_0 .

In the prefix before entry, delete path segments between repeated pairs consisting of an outside state and elapsed time modulo d . Each deletion preserves the endpoint and length modulo d . The resulting path of length $h \equiv h_0 \pmod{d}$ has distinct outside product vertices at all positions before entry. Shorten the suffix analogously to length $\beta \equiv \beta_0 \pmod{d}$, with distinct outside product vertices at all positions after exit.

The outside states occurring before entry and after exit in the original walk are disjoint. A state occurring in both could reach C and be reached from C , contradicting its being outside that SCC. The shortened paths use only these original states. If their outside-state counts are a and b , then $h \leq ad$, $\beta \leq bd$, and $a + b \leq s - c$, proving (93). This also covers zero-length access or exit paths.

The interior segment of the original walk has length congruent to $\kappa(v) - \kappa(u)$ modulo d . Because $d \mid p$, the two shortenings yield (94). The original entry pair lies in the stripe by the cyclic-class invariant along the

path to the known stripe vertex. Preserving entry time modulo d therefore keeps $(h \bmod p, u)$ in the same *full* stripe, despite a possible change of its phase modulo p . $\square$

[Added 1 October 2026, batch 73: The proof uses no determinism, only that the stripe is a full stripe over an original SCC. Part VII states the lemma for the universal stripe of an NFA as Lemma 79.2 and proves it by this argument, which is printed only here.]

Proof of Theorem 73.1. First assume $s \geq 2$. For the component C , Schwarz's theorem supplies an internal relation-power transient

$$1 \leq k \leq (c - 1)^2 + 1.$$

For every length $\ell \geq k$, an internal walk from x to y exists exactly when $\ell \equiv \kappa(y) - \kappa(x) \pmod{d}$. To justify using the numerical threshold rather than an unspecified asymptotic one, enlarge its matrix period to a multiple of d and propagate the equality of matrix entries forward from ℓ to the sufficiently long regime of the cyclic-class path criterion. The stable value is already attained at ℓ .

Consider a nonempty hull word w of length $n \geq t$, so its residue r is live. Use the stripe and entry/exit paths above and set

$$A = \max(h, s), \quad B = \max(\beta + k, s).$$

If $n \leq A + B$, singleton fragments already give rank at most $A + B$. Otherwise, pad the h -step access path by an arbitrary internal C -walk of length $A - h$. Such a walk exists for every nonnegative length because each vertex of the cyclic strongly connected graph has an internal outgoing edge. The padded path reaches a stripe vertex at phase A .

Read the target central interval $w[A : n - B]$ from that vertex. Every letter is stable-permitted by Lemma 73.2: its prefix length is at least $A \geq s$, and even its last letter has suffix length at least $B \geq s$. Stripe safety keeps the deterministic run inside C . Let z be its endpoint. Its class is

$$\kappa(z) \equiv \kappa(u) + (A - h) + (n - A - B) = \kappa(u) + n - h - B \pmod{d}.$$

Using (94) gives

$$\kappa(v) - \kappa(z) \equiv B - \beta \pmod{d}.$$

Because $B - \beta \geq k$, there is an internal completion of exactly $B - \beta$ steps to v . Append the fixed β -step exit path. The resulting accepting seed has total length exactly n and agrees with the target on its entire central interval. Hence that interval is one aligned admissible fragment, and

$$\begin{aligned} \rho_L(w) &\leq A + B + 1 \\ &\leq h + \beta + k + 2s + 1 \\ &\leq (s - c)d + (c - 1)^2 + 2s + 2 \\ &\leq (s - c)c + (c - 1)^2 + 2s + 2 \\ &= c(s - 2) + 2s + 3 \leq s^2 + 3. \end{aligned}$$

For $n < t$, the singleton bound is at most $n < t \leq (s - 1)^2 + 1$. Empty slices require no words to be bounded, and the empty-word convention contributes at most one. Finally a complete one-state DFA accepts either all words or none, so $R(L) = 1$ and the asserted estimate also holds for $s = 1$. $\square$

Combining Theorem 73.1 with the deterministic lower family gives

$$\frac{(s - 7)^2 + 2}{s^2} \leq \frac{F_{\text{DFA}}(s)}{s^2} \leq 1 + \frac{3}{s^2} \quad (s \geq 8),$$

and therefore the sharp leading constant one. The result does not determine the optimal lower-order terms or exact extremal function. Determinism enters through the safe-core characterization and its strict survivor-count decrease; nondeterministic images can expand, so this proof does not supply a leading-one NFA bound. [Added 1 October 2026, batch 73: Part VII supplies one by another route, robust compression in a primitive component (Theorem 77.2) in place of the safe-core step; it uses Lemmas 73.2 and 73.3 in their NFA forms.]

74 Finite verification and attribution

The package contains four Python standard-library checkers and their saved outputs. They have different scopes.

- `check_stripes_exhaustive.py` finds the exact eventual cycle of (P_i, R_i) , enlarges its period to include SCC periods, computes the safe core by repeated deletion, and compares it with the stripe criterion. It checks every labelled complete binary DFA with fixed initial state and every final set on one, two, and three states: $2 + 64 + 5,832 = 5,898$ automata and 26,881 live residues. It also checks uniform stripe viability.
- `check_stripes_random.py` checks 1,519 seeded random complete binary DFAs on one to eight states and all their live residues. Its random seed is 153735.
- `check_rank_family.py` verifies 24 path-based alternating rank witnesses for $2 \leq m \leq 25$, reaching rank 578. It also literally tests 14,308 seed words for $2 \leq m \leq 8$ and $2 \leq n \leq 10$; these are seed-membership tests, not 14,308 separate rank computations.
- `check_gate_completion.py` independently verifies 1,829 gate-state completion cases over $2 \leq m \leq 60$, the missing return at g , and 19 alternating rank witnesses for $2 \leq m \leq 20$.

Run the scripts with assertions enabled, without Python's `-O` option. The tests corroborate the finite constructions and predicates. They do not prove the universal stripe lemmas, the complexity classification, correctness of every CRT reduction instance, or the quadratic asymptotic theorem. The proofs supply those arguments. The new uniform $s^2 + 3$ bound is an analytic deduction from the joint path budget; no finite test here exhausts that universal claim. No Lean or other proof-assistant formalization is included.

The contribution claimed relative to the pinned source is the SCC-stripe certificate and its `coNP` consequence, the crossover wrapper around classical modular testers, the complete-DFA quadratic-rank construction, and the joint entry/exit argument giving its sharp leading constant. The rank clocks, periodic interiors, deterministic safe cores, CRT arithmetic, selector-to-cycle devices, and relation-power bound are credited background. A comprehensive search excluding all equivalent earlier crossover results has not been completed, and no global first-discovery claim is made.

Open refinements include the following. [write, batch 73] The manuscript lists them in one sentence; they are printed as numbered questions, in the same order and words; their bracketed titles were added.

Question 74.1 (Exact DFA extremum). What is the exact $F_{\text{DFA}}(s)$?

Question 74.2 (Lower-order terms). What are the optimal lower-order terms beyond s^2 ?

Question 74.3 (The deterministic family). What is the exact rank of the displayed deterministic family?

Question 74.4 (One automaton). What is the complexity of computing an individual DFA's exact finite rank or least stabilization generation?

The sharp leading DFA rank constant and the finite-stabilization decision classification are settled by the results above, rather than left on that list. [write, batch 73] The window of Theorem 67.2 has width $s^2 + 3 - ((s - 7)^2 + 2) = 14s - 48$; Question 74.1 refines Part IV's Question 53.1 for complete DFAs, and Question 74.3 asks where in $[g + 2, g + 4]$ the rank of L_m^{dfa} lies. No statement of this Part is formalized; Part II's Question 24.8 applies to it as well.

[write, batch 73] *Shipped layout.* The paragraphs above describe the delivered package. In this report its files are shipped, byte-identical, as follows; the manuscript and its PDF are printed or replaced by this report, the delivery README and the revision note REVISION.md by the report's README, and the checksum list SHA256SUMS was verified and retired.

Shipped path	Delivered as
code/06-dfa-crossover-check_stripes_exhaustive.py	check_stripes_exhaustive.py
data/06-dfa-crossover-stripes-exhaustive.json	stripes-exhaustive.json
code/06-dfa-crossover-check_stripes_random.py	check_stripes_random.py
data/06-dfa-crossover-stripes-random.txt	stripes-random.txt
code/06-dfa-crossover-check_rank_family.py	check_rank_family.py
data/06-dfa-crossover-rank-family.json	rank-family.json
code/06-dfa-crossover-check_gate_completion.py	check_gate_completion.py
data/06-dfa-crossover-gate-completion.json	gate-completion.json
code/06-dfa-crossover-build.sh	build.sh

The four checkers print their records to standard output and write no file; the report's README gives commands that keep the output away from the shipped records. When the batch was placed they were rerun on a copy with Python 3.14.4. Three took between seven and eleven seconds each and printed the bytes of their shipped records apart from CRLF line endings. The random stripe checker needs about six minutes on that machine, the two eight-state cases alone about 72 and 149 seconds; it was run in three parts with the same random stream, and all 1,519 automata passed. The build script changes to its own directory and runs `pdflatex` twice on `dfa-crossover-stabilization.tex`, the delivered manuscript, which is printed here as Part VI and is not shipped; it must not be run in `code/`.

Part VII. Sharp leading constant for constrained crossover rank

[write, batch 73] *Part VII was added on 1 October 2026 in batch 73, from one manuscript written against Parts I–III as printed above; it knew Parts IV and V only as separate manuscripts and Part VI not at all. It answers Part IV’s Question 53.1 as to the leading constant: for binary NFAs $(s - 3)^2 \leq F(s) \leq s^2 + 8s$, so $F(s)/s^2 \rightarrow 1$, and for a q -letter alphabet the upper bound is $s^2 + \max\{8, 2q - 4\}s$. It computes $F(1), \dots, F(4) = 1, 2, 5, 7$, the last three with computer assistance, and gives an exact algorithm and a witness cutoff for the maximal rank of one automaton. The exact $F(s)$ for $s \geq 5$ stays open. Section 75 records the provenance and fixes the notation of the addition.*

75 [write, batch 73] Provenance and notation of Part VII

75.1 One manuscript, one addition

Part VII is batch-73 manuscript 21, *Sharp Leading Constant for Constrained Crossover Rank*, dated 1 October 2026 (31 US-letter pages as delivered). It is an AI-assisted research note: its title page and PDF metadata read “Research note prepared with OpenAI for Vladimir Reshetnikov”, quoted here and in the report’s README only. It arrived as the archive `common-path-rank-package.zip` in ProveIt commit `f8c3a392a` and was placed in commit `5e4f1263c`. Its source files are named `hamiltonian-rank.tex` and `.pdf`, a name its README says was kept from an earlier delivery; that delivery is not in the repository. Its bibliography and its file `source-pins.json` pin ProveIt commit `1bd730779` (in full `1bd730779e8bc44c8976dcd9f38cad86a0d21c26`) for this report, where `article.tex` held Parts I–III as printed above (blob `ff492c83`), and commit `aee32ad45` for Part IV’s manuscript, then the archive `docs/incoming/quadratic-finite-rank-package.zip`; they identify Part V’s manuscript by the SHA-256 of its source file, which is the source of Part V. So “the pinned report” is this report and its “Part II” is Part II, and “the companion notes” are Parts IV and V.

The earlier version. The same batch delivered an earlier version of this manuscript, batch-73 manuscript 24 (archive `common-path-rank-package (1).zip`, 19 pages, same pins). The review receipt of manuscript 21 names it as the previous version by the SHA-256 of its source. Every result and proof of it is in manuscript 21 verbatim, apart from the name F for F_2 ; manuscript 21 adds the manuscript’s Sections 6–9 (Sections 81–84 here: the finite-alphabet bound, the linear family, the exact algorithm and the censuses). The one sentence of the earlier version that manuscript 21 withdraws, that the same leading coefficient for each fixed alphabet of at least three letters “also remains open here”, is answered by Theorem 81.1. It was superseded, is not printed, and nothing of it is shipped; it survives in the arrival commit.

Its checkers, census programs, receipts and logs and its build and reproduction scripts are shipped beside the files of Parts I–VI with the prefix `07-sharp-constant-` (45 files; the delivered subdirectories are flattened into the name, so `exact/four_state/audit_four.cpp` is `code/07-sharp-constant-exact-four_state-audit_four.cpp`). Two delivered files are not shipped again, because they are byte-identical to Part IV’s shipped checker and record. Its manuscript, delivery README and PDF are not shipped and survive in the arrival commit, and its checksum list `SHA256SUMS` (50 files, all verified at placement) was retired by the repository’s policy. Section 88 and the report’s README list the shipped names.

Every result, proof, example, table, limitation and remark of the manuscript is printed here, in its order,

except four proofs that repeat proofs printed in earlier Parts and are replaced by pointers: those of Lemma 78.1 and Corollary 83.4 (Part II), of Lemmas 79.1 and 79.2 (Part VI, which states them for DFAs), and of Proposition 80.1 (Part IV). The statements are kept. Its proof of Lemma 83.1, Part V’s lemma, is kept as a second proof. Its title page and abstract became this section; its Sections 1–13 are Sections 76–88 (manuscript Section n is Section $n + 75$; its main theorems are Theorems 76.2 and 81.1); it has no appendix and no numbered questions. Text marked [write, batch 73] was added when the manuscript was written into this report. No theorem was strengthened or weakened.

What Part VII adds, and what it repeats.

- *Answered:* Part IV’s Question 53.1 (“What is the best leading constant, or the exact function $F(s)$?”), as to the leading constant: it is one, $(s - 3)^2 \leq F(s) \leq s^2 + 8s$ for $s \geq 4$ (Theorem 76.2); and as to the exact function for $s \leq 4$: $F(1), \dots, F(4) = 1, 2, 5, 7$ (Theorem 84.1, computer-assisted for $s = 2, 3, 4$). This also settles the constant left in Part II’s Question 24.2 and narrows Part V’s $[1, 2]$ to 1.
- *New:* robust compression in a primitive factor-universal component (Theorem 77.2, with Lemmas 77.3–77.5); a universal strongly connected component and its full stripe (Lemmas 78.2 and 78.3); the imprimitive and one-symbol cases; the finite-alphabet bound $s^2 + \max\{8, 2q - 4\}s$ with its common-anchor lemmas (Theorem 81.1, Lemmas 81.2–81.4) and an all-symbol lower family; the linear family with $R = 2s - 1 = 2e + 1$ and its slice maxima (Proposition 82.1); the exact algorithm and the witness cutoff (Theorems 83.2 and 83.3); the small extrema (Theorem 84.1); the Hamiltonian and near-Hamiltonian bounds and the barrier family (Theorems 85.1 and 86.1, Proposition 87.1).
- *Shared with Part VI:* Lemmas 79.1 and 79.2 are Part VI’s Lemmas 73.2 and 73.3, here for NFAs (and with a liveness clause), with the same proofs, which are printed once, in Part VI. The two manuscripts were written separately and do not cite each other.
- *Restated:* Lemma 78.1 and Corollary 83.4 are Part II’s Theorem 18.3 (necessity, via Theorem 18.4) and its bound (42); Proposition 80.1 is Part IV’s Theorem 47.1, and the manuscript says so; Lemma 83.1, the parser after it and the periodic-middle graph are Part V’s (second proof kept).
- *Imported and credited:* the primitive-exponent results of Kirkland, Olesky and van den Driessche [30] (with Lewin and Vitek’s two-cycle-length result, which they credit) and the classical bound $e \leq (n - 1)^2 + 1$.
- *Not affected:* the exact $F(s)$ for $s \geq 5$, the optimal additive term, and a leading constant uniform in the alphabet size when q is comparable to s are not determined; Part IV’s Questions 53.3 and 53.4, Part V’s and Part VI’s questions, and the questions of Parts I–III are untouched. For complete DFAs Part VI’s bound $s^2 + 3$ is sharper than $s^2 + 8s$.

75.2 The manuscript’s summary and status

Abstract, as delivered. For every s -state ε -free NFA over a q -letter alphabet, $q \geq 2$, with finite maximal aligned-fragment rank, we prove $R(L) \leq s^2 + \max\{8, 2q - 4\}s$. For every $s \geq 4$, an explicit family using every alphabet symbol has $R(L) = (s - 3)^2$. Hence $F_q(s)/s^2 \rightarrow 1$ for each fixed alphabet and even when $q = o(s)$. The proof combines primitive-graph compression, a circulation-based common-anchor theorem for selected monochromatic cycles, and a reduction through periodic viable layers. We also give an exact weighted-parser algorithm and a complete witness-length cutoff. Independent exhaustive computations establish the binary values $F_2(1) = 1$, $F_2(2) = 2$, $F_2(3) = 5$, and $F_2(4) = 7$, with a proved transient pruning rule in the four-state case. For every $s \geq 3$, a uniform s -state family has slice maxima $\min(N, 2s - 1)$ for $N \geq 1$ and attains the exponent bound $R = 2e + 1$ with $e = s - 1$. Sharper Hamiltonian bounds and a quadratic-exponent,

linear-rank barrier family are retained. Ordinary all-size proofs are distinguished from the computer-assisted censuses and finite consistency checks. No exact formula for $F_2(s)$ at $s \geq 5$, unrestricted alphabet-growth leading coefficient, proof-assistant certification, or literature-wide priority is claimed.

Status, as delivered. The manuscript has no separate status box; its abstract ends with the disclaimer above, and Section 88 states its scope and the three roles of its evidence. Its delivery README adds: no alphabet-size-independent additive term or leading-one conclusion for unrestricted alphabet growth is claimed; the article keeps $R(\emptyset) = 1$, whereas the census programs record an empty hull as raw value 0, a difference that affects no extremal value; the four-state counts are pruned representative counts, not a full four-state rank distribution; the bounded tests do not substitute for an all-size proof; this is an unrefereed note with ordinary and computer-assisted proofs, not a proof-assistant certificate; no literature-wide priority claim is made. It also says that every page of the delivered PDF was visually inspected, and that raw million-row listings and compiled executables were omitted from the package.

75.3 Notation beside Parts I–VI

Part VII uses the report's rank ρ_L , hull $\text{Rec}(L)$, supremum rank $R(L)$, live residues, phase-admissible words and stable layers, and Part II's slice maximum $M_L(N)$. One symbol was renamed when the manuscript was written in: its backward exact-length layers B_j are printed R_j , Part II's name for the same sets, because B_s is Part II's constant $(2s + 5)^2$ and the manuscript itself uses B for three other things. No normalization changed. In the manuscript, letters are indexed from one in the definition of the hull and boundaries from zero, as in Part I.

In Part VII	Meaning in Part VII	Not to be confused with
$F_2(s), F_q(s)$	largest finite $R(L)$ over NFAs with at most s states on 2, or q , letters; the subscript is the alphabet size	Part II's $F(s)$, which is $F_2(s)$; $F_3(s)$ is ternary, not $F_2(3) = 5$; Part VI's $F_{\text{DFA}}(s)$
R_j (manuscript B_j)	backward exact-length layers	Part II's constant B_s ; the B of Sections 77, 79 and 81
$\Gamma_{r,a}$	stable alphabet, Part II's $C_{r,a}$	the alphabet $\Gamma = \{0, 1, @\}$ of Parts I and II; Part III's graph Γ_R
C, c, d, κ	the original component under a universal stripe, its size, period and cyclic classes	C_q ; $C_{r,a}$; Part V's index κ
$C_q = \max\{8, 2q - 4\}$	coefficient of the linear term	the stable alphabets $C_{r,a}$ of Parts II, V, VI
$\omega_n = (n - 1)^2 + 1$	the classical bound on primitive exponents	Part III's clique number $\omega(\Gamma_R)$
e, e_0, e_C ; $j < k$	exponents; the two cycle lengths	
m	$ V(U) $ in Sections 77 and 81; a cyclic class size in Section 79; the gate size $s - 2$ in Section 80	Part II's $m = s + 2$; Part IV's gate size $m = s - 2$
T, V, H	imprimitive length threshold; vertex bound $p(2^s - 1)$; witness cutoff	the operator $T(L)$, Part II's T_s ; Part IV's H_m ; Part V's N
A_r, D_r	boundary operators of the exact algorithm	Part II's $A_{n,i}$ and $D_L(n)$
$\delta(S, x)$	subset image under an NFA	Part VI's DFA transition δ
H_m, T_m, L_m	in the manuscript's proof of Proposition 80.1, not reprinted: the semigroup, the gate lengths, the gate language	<i>swapped</i> against Part IV's S_m, H_m ; Part IV's L_m^{gate}

The manuscript's macros for ε , ρ and Rec are the report's; its unused macro $\backslash\text{ex}$ was not carried over. Its `lmodern` and `amssymb` packages are not loaded.

76 Aligned fragments and the main results

[write, batch 73] “The preceding reports” and “the pinned report” of this Part are this report as it stood at the pin, Parts I–III; “the companion notes” are the manuscripts of Parts IV and V. The manuscript’s citations of all three are replaced by internal references. Its $F_2(s)$ is Part II’s $F(s)$, and its R_j (written B_j in the manuscript) are Part II’s backward layers; Section 75.3 lists the other symbols.

Let Σ be a nonempty finite alphabet. The definitions below apply to every such alphabet; Sections 77–79 first establish the binary upper bound, with $\Sigma = \{0, 1\}$. For $L \subseteq \Sigma^*$ write $L_N = L \cap \Sigma^N$. Boundaries are indexed from zero: $w[u : v]$ consists of the letters in positions $u + 1, \dots, v$.

Definition 76.1. If $L_N \neq \emptyset$ and $N \geq 1$, the *coordinate hull* at length N is

$$\text{Rec}(L)_N = \prod_{i=1}^N \{x_i : x \in L_N\}.$$

It is empty when L_N is empty, and $\text{Rec}(L)_0 = L_0$. Put $\text{Rec}(L) = \bigcup_{N \geq 0} \text{Rec}(L)_N$. For a target $w \in \Sigma^N$, a nonempty interval $[u, v]$ is *admissible* if a full seed $x \in L_N$ satisfies $x[u : v] = w[u : v]$. Its *aligned-fragment rank* $\rho_L(w)$ is the least number of consecutive admissible intervals partitioning $[0, N]$, or ∞ if none exist. Set $\rho_L(\varepsilon) = 1$ if $\varepsilon \in L$, and ∞ otherwise. Finally define

$$R(L) = \sup(\{1\} \cup \{\rho_L(w) : w \in \text{Rec}(L)\}).$$

Every fragment has its original position and a seed of the same full length. Different fragments may use different seeds. In particular, finite rank is equivalent to hull membership, and $\rho_L(w) \leq N$ for every nonempty hull word of length N . The supremum defining $R(L)$ excludes words outside the hull; its value can still be infinite if the finite ranks are unbounded. These are the fragment conventions used for constrained crossover in the preceding reports (Part I’s Definitions 2.4 and 3.1, Part II’s Definition 15.1 and Corollary 15.3, and Part V’s Section 56).

An ε -free NFA here consists of a state set Q , transition relations $\Delta_a \subseteq Q \times Q$ for $a \in \Sigma$, and sets $I, F \subseteq Q$. Its underlying digraph has edge relation $E = \bigcup_{a \in \Sigma} \Delta_a$. An edge may carry several letters. For $S \subseteq Q$ and $x \in \Sigma^*$, let $\delta(S, x)$ be the possible endpoints of paths labeled x from S . The NFA is *factor-universal* if

$$\delta(Q, x) \neq \emptyset \quad \text{for every } x \in \Sigma^*.$$

Thus all words label some walk, with both endpoints unrestricted. This is a condition on the labeled graph, not a claim that $L = \Sigma^*$ for the chosen I, F .

A digraph is *primitive* if some positive length admits a walk between every ordered pair of vertices; its least such length is its *exponent* e . Equivalently, it is strongly connected and its cycle lengths have gcd one. A Hamilton cycle visits all vertices once.

Theorem 76.2 (General binary NFA bound). *Let an s -state, ε -free binary NFA have arbitrary initial and final sets, where $s \geq 1$. If $R(L) < \infty$, then*

$$R(L) \leq s^2 + 8s.$$

Consequently, if $F_2(s)$ is the supremum of finite values $R(L)$ over binary NFAs with at most s states, then

$$(s - 3)^2 \leq F_2(s) \leq s^2 + 8s \quad (s \geq 4), \quad \boxed{\lim_{s \rightarrow \infty} \frac{F_2(s)}{s^2} = 1.} \quad (95)$$

For alphabet size $q \geq 2$, write $F_q(s)$ for the supremum of finite ranks recognized by NFAs over a fixed q -letter alphabet with at most s states. The subscript always records alphabet size and the argument records state count. The same extremal convention results from using exactly s states, since unreachable states may be added. The main generalization, proved in Section 81, is

$$(s - 3)^2 \leq F_q(s) \leq s^2 + \max\{8, 2q - 4\}s \quad (q \geq 2, s \geq 4). \quad (96)$$

Thus the leading coefficient is one for each fixed finite alphabet and for every alphabet-size sequence $q = o(s)$. In particular $s^2 + 8s$ works for every alphabet of size two through six. The exact binary small-state results, established separately in Section 84, are 1, 2, 5, 7 for one through four states. The subscript in $F_q(s)$ must not be confused with the state argument: the ternary function is $F_3(s)$, while the binary three-state value is $F_2(3) = 5$.

Sections 77–79 first prove the binary upper bound, using two monochromatic cycles and an alphabet-independent phase-component reduction. Section 80 proves the quadratic lower construction. Section 81 supplies the small-family common-anchor theorem and extends the bound to every finite alphabet. Section 82 gives the sharp linear family and its exact slice profile. Sections 83–84 separate the exact algorithm and pruning proofs from the exhaustive binary computations. Section 85 preserves sharper bounds for Hamiltonian and near-Hamiltonian graphs, followed by the barrier to a total-state-subset approach. None of the arguments assumes determinism, a single initial state, or a total labeled state subset.

The leading-coefficient result sharpens the earlier binary comparison $(s - 3)^2 \leq F_2(s) \leq 2(s - 1)^2 + 3$ from the companion notes (Part IV’s Theorem 47.1 and Part V’s Corollary 64.2). The proof here is complete apart from the explicitly cited classical primitive-digraph exponent results. The periodic-interior implication needed for the reduction is reproved, and no global transient estimate enters the leading-coefficient bound.

[write, batch 73] Theorem 76.2 answers Part IV’s Question 53.1 as to the leading constant, which is one, for NFAs over any fixed finite alphabet. Two consequences are the write’s observations, not the manuscript’s. First, the bound $s^2 + 8s$ needs no assumption on the initial states, so with Part IV’s single-initial-state family (Section 52) it gives $(s - 4)^2 \leq F_{\text{one}}(s) \leq s^2 + 8s$ and $F_{\text{one}}(s)/s^2 \rightarrow 1$ as well. Second, it does not make Part V’s bound obsolete for small s : $2(s - 1)^2 + 3 < s^2 + 8s$ exactly when $s \leq 11$ (for example $203 < 209$ at $s = 11$, and $240 < 245$ the other way at $s = 12$). For complete DFAs, Part VI’s $s^2 + 3$ (Theorem 73.1) is smaller than both.

77 Robust compression inside a primitive component

The general reduction needs more than a rank theorem for a component’s own coordinate hull. An interval taken from a hull word of the whole NFA need not belong to that smaller hull. The following formulation explicitly permits its singleton witnesses to come from outside the component.

Definition 77.1. Fix vertices u, v in a labeled graph and a target x of length $M \geq 1$. A *permissive partition* is a consecutive partition of x in which each piece is either a singleton with no internal witness required, or agrees on the same interval with the label of a complete M -edge walk from u to v . The endpoints and full length of these witnesses are fixed; different pieces may use different walks.

Theorem 77.2 (Robust primitive bound). *Let an n -vertex binary labeled graph be primitive and factor-universal. For every u, v , every $M \geq 1$, and every binary target of length M , there is a permissive partition with at most*

$$n^2 + 8n$$

pieces. Consequently any NFA with this graph and nonempty initial and final sets satisfies $R(L) \leq n^2 + 8n$.

The consequence follows by choosing $u \in I, v \in F$: replace every licensed singleton by its full-language hull witness. The assertion does not claim that every short length admits an internal u -to- v walk. At such lengths an all-singleton partition is allowed.

77.1 Exponent reduction and the two-cycle geometry

For a primitive graph of exponent e , a target of length $M > 2e$ has a single internally witnessed middle interval $[e, M - e]$. Indeed, factor universality gives a path labeled by that interval, and e -step walks from u and to v complete it to length M . The two boundaries can be licensed singletons. Hence the permissive rank is at most $2e + 1$; when $M \leq 2e$, use only singletons.

Put $\omega_n = (n-1)^2 + 1$. The classical results of Kirkland, Olesky and van den Driessche [30, Theorems 1.1–1.2 and the proof of Theorem 1.2] give the following input. If $n \geq 6$ and

$$e \geq \lfloor \omega_n/2 \rfloor + 2, \tag{97}$$

then the only elementary cycle lengths are two coprime integers $j < k \leq n$, and

$$e \leq n + j(k - 2). \tag{98}$$

Below this threshold, $2e + 1 \leq (n - 1)^2 + 4$. If $j \leq n/2$, equation (98) gives $e \leq n^2/2$. For $2 \leq n \leq 5$, the classical bound $e \leq \omega_n$ suffices, and the one-state graph has both loop labels. All these cases satisfy Theorem 77.2. It remains to treat

$$n \geq 6, \quad n/2 < j < k \leq n, \quad \gcd(j, k) = 1, \quad e \leq jk. \tag{99}$$

The final inequality follows from $n - 2j < 0$ in (98).

Lemma 77.3 (Order of cycle intersections). *If a digraph has only the elementary cycle lengths $j < k$, where $j \geq 3$ and $\gcd(j, k) = 1$, any two elementary cycles encounter their common vertices in the same cyclic order.*

Proof. Otherwise there are common vertices a, b, c ordered a, b, c on one cycle and a, c, b on the other. The path from a to b on the first cycle followed by the path from b to a on the second is a nonempty closed walk. Do the same for b to c and for c to a . These three nonempty closed walks partition the two original cycles' edge multiset, with multiplicity. Decomposing them into elementary cycles produces at least three cycles whose total length is $2j, j + k$, or $2k$.

Each such total has a unique nonnegative representation $Aj + Bk$, namely its original representation $A_0j + B_0k$ with $A_0 + B_0 = 2$. To see this, any other representation has $B = B_0 + tj$ and $A = A_0 - tk$ for an integer t . Since $A_0, B_0 \leq 2 < j < k$, neither sign of nonzero t preserves nonnegativity. Every decomposition must therefore contain exactly two elementary cycles, a contradiction. With fewer than three common vertices the order assertion is automatic. $\square$

Lemma 77.4 (A circular chain of alternatives). *Suppose two elementary cycles meet and satisfy Lemma 77.3. Their union is a circular chain of pairs of directed paths between consecutive common vertices, allowing internally simple return paths to the same vertex when there is just one common vertex. Each pair either is a single shared edge or consists of two internally disjoint paths, called a bubble. Choices of paths in different bubbles are independent and produce elementary cycles. If the two original cycles have equal length, every bubble has equal-length alternatives. If their lengths are $j < k$, exactly one bubble has unequal alternatives, of lengths differing by $k - j$.*

Proof. Between consecutive common vertices, the two paths have no further common internal vertex; distinct pairs have disjoint interiors. Thus any choice visits each common vertex once and is an elementary cycle. If both original cycles have length j , replacing any one path cannot shorten a cycle below j , so every difference, measured from the first cycle to the second, is nonnegative. Their sum is zero, so each difference is zero. For two k -cycles use the analogous prohibition against exceeding k .

For cycles of lengths $j < k$, every difference measured from the j -cycle is nonnegative. A positive difference must be exactly $k - j$, because otherwise a single replacement would produce a cycle of a forbidden intermediate length. The differences sum to $k - j$, proving uniqueness of the unequal bubble. The one-common-vertex case consists of a single bubble and obeys the same argument. $\square$

Factor universality supplies a monochromatic elementary cycle C_0 for 0 and one C_1 for 1: read a constant word of length at least n and extract a cycle. Under (99), these cycles meet, since each has more than $n/2$ vertices. Write $U = C_0 \cup C_1$ and $m = |V(U)| \leq n$. Every shared edge carries both labels.

Lemma 77.5 (Equal-bubble fragments). *Inside a complete walk that traverses the bubbles of U , a length- t equal bubble requires cuts only where consecutive target letters differ within that bubble. No cut is required just because the target crosses a common vertex or a shared edge. Give charge zero to shared edges and to the first edge of each equal bubble path, and charge one to its other edges. A segment meeting only equal-length bubbles and shared edges, of charge W and with every touched bubble in a complete ambient traversal, can be covered by at most $W + 1$ internally witnessed fragments.*

Proof. After cutting at each internal symbol change, every fragment's portion in a given bubble is constant. Choose its 0- or 1-monochromatic branch accordingly. Choices for different bubbles are independent, and each preserves both endpoints and length. Shared edges permit both symbols. Thus a fragment can cross a bubble endpoint and continue into further bubbles without another cut.

For a fragment starting or ending inside a bubble, replace the whole containing traversal in its complete witness walk, including the portion outside the fragment. The witness still has the original total length and endpoints. Each possible internal change is charged to the second of its two edges, so the number of cuts is at most W . Different fragments may use different branch choices. $\square$

All applications below use a repeated-cycle block based at a common vertex of C_0, C_1 . Its bubbles are therefore complete eligible traversals, even when the segment being covered cuts through one of them.

77.2 Equal lengths of the monochromatic cycles

Suppose $|C_0| = |C_1| = \ell \in \{j, k\}$ and write r for the other length. Set $h = m - \ell$. If the equal bubbles have lengths $t_1, \dots, t_b$, vertex counting gives

$$h = \sum_{i=1}^b (t_i - 1).$$

Every elementary cycle of U has length ℓ and charge h under Lemma 77.5. Put $d_0 = h/\ell$, so $0 \leq d_0 < 1$, since $m \leq n < 2\ell$. Successively deleting elementary cycles from a U -walk segment of length x leaves a simple path of at most $m - 1$ edges, each of charge at most one. Its total charge is consequently at most

$$d_0x + (1 - d_0)(m - 1). \quad (100)$$

Cycle deletion only counts charge; fragments stay at the original positions.

Choose an elementary r -cycle J . It meets C_0 because $r + \ell > n$. Let $z \in J \cap C_0$ and let $y \in C_0 \cap C_1$. Choose a path from u to z of length $a \leq n - 1$, the forward C_0 path from z to y of length $t \leq \ell - 1$, and a path from y to v of length $b \leq n - 1$. Their combined length is

$$D_0 = a + t + b \leq 2n + \ell - 3.$$

For $M - D_0 \geq (\ell - 1)(r - 1)$, the elementary two-generator Frobenius formula yields

$$M - D_0 = \alpha r + \beta \ell, \quad \alpha, \beta \geq 0.$$

Reducing α modulo ℓ increases β and gives $0 \leq \alpha < \ell$. Use one complete M -edge walk, in this order:

$$u \rightarrow z; \quad J^\alpha; \quad z \rightarrow y \text{ along } C_0; \quad C_0^\beta \text{ at } y; \quad y \rightarrow v.$$

Its single J block has length $K_0 = \alpha r < jk$. Replacements inside the repeated- C_0 block are legitimate complete-walk witnesses.

If the representation threshold fails, the singleton bound is enough, because

$$M < D_0 + (\ell - 1)(r - 1) \leq jk + 2n - r - 2 < n^2 + 2n.$$

Otherwise cover the whole target if $M \leq 2e$, or only the two windows $[0, e)$ and $[M - e, M)$ if $M > 2e$, supplementing them with the universal middle fragment. Let $T \leq 2e$ be the number of covered window positions and K the number lying in the single J block. Importantly, $K \leq K_0 < jk$ in total across both windows. The fixed connecting paths cost at most D_0 licensed singletons. The repeated- C_0 portions form at most two segments. Equations (100) and (99) give a total of at most

$$\begin{aligned} d_0T + (1 - d_0)K + D_0 + 2(1 - d_0)(m - 1) + 3 &\leq 2d_0e + (1 - d_0)jk + D_0 + 2m + 1 \\ &\leq jk(1 + d_0) + D_0 + 2m + 1 \\ &= rm + D_0 + 2m + 1 \\ &\leq rn + 4n + \ell - 2 \leq n^2 + 5n. \end{aligned}$$

The constant three allows one extra piece for each of the at most two repeated-cycle segments and one middle piece. If the monochromatic cycles coincide, $h = 0$ and every edge is doubly labeled; the argument remains valid.

77.3 Unequal lengths of the monochromatic cycles

After interchanging symbols if necessary, take $|C_0| = j$ and $|C_1| = k$. Let the unique unequal bubble have path lengths p_0 and $q_0 = p_0 + k - j$. If B is the sum of $t - 1$ over the equal bubbles, then

$$m = j + B + q_0 - 1.$$

Use the equal-bubble charges already defined, but charge every edge of the unequal bubble by one, with an additional unit on its first edge. Thus edge charges belong to $\{0, 1, 2\}$. A length- a traversal of the unequal bubble has charge $a + 1$, enough to cut before it and after every position and to use licensed singleton pieces there. Cuts outside the covered segment are unnecessary. In the other bubbles, Lemma 77.5 applies. Any segment has a permissive partition of size at most its charge plus one, also when it starts or ends inside the unequal bubble.

The charges of the j - and k -cycles are respectively

$$B + p_0 + 1 = m - k + 2, \quad B + q_0 + 1 = m - j + 2.$$

Since $j + k \geq n + 2 \geq m + 2$, their maximum mean charge is

$$\mu = \frac{m - j + 2}{k} \leq 1,$$

for

$$j(m - j + 2) - k(m - k + 2) = (k - j)(j + k - m - 2) \geq 0.$$

All cycles in U arise by branch choices, so these are the only possible means. Deleting cycles leaves a simple path of at most $m - 1$ edges, each of charge at most two. A segment of length x therefore has charge at most

$$\mu x + (2 - \mu)(m - 1).$$

Choose a common vertex y and fixed paths $u \rightarrow y$ and $y \rightarrow v$ with total length $D_0 \leq 2n - 2$. If $M - D_0 \geq (j - 1)(k - 1)$, write $M - D_0 = \alpha j + \beta k$ and use the complete walk consisting of the entry path, α copies of C_0 and β copies of C_1 based at y , and the exit path. Below this threshold the singleton bound is less than $n^2 + 2n$.

As before, cover the full target when $M \leq 2e$, or the two boundary windows otherwise, with a separate middle fragment. There are at most two covered U segments. The number of pieces is at most

$$\begin{aligned} 2\mu e + D_0 + 2(2 - \mu)(m - 1) + 3 &\leq 2j(m - j + 2) + 6n - 3 \\ &\leq 2j(n + 2 - j) + 6n - 3 \\ &\leq \frac{(n + 2)^2}{2} + 6n - 3 \\ &= \frac{n^2}{2} + 8n - 1 \leq n^2 + 8n. \end{aligned}$$

Together with the preliminary cases, this proves Theorem 77.2. In particular, every nonsingleton witness has exactly the required M edges from the fixed u to the fixed v . No target-hull assumption for the component entered the proof.

78 Stable phases and a universal cyclic component

We supply the finite-rank implication of the periodic-interior criterion from Part II (Theorem 18.3). Only an eventual period is needed here; its size will never be charged to rank.

For the underlying relation E , use the exact-length layers

$$P_i = E^i(I), \quad R_j = \{q : \text{a } j\text{-edge walk leads from } q \text{ to } F\}.$$

For a target of length N , a word x can occur as an aligned seed fragment on $[i, j)$ exactly when

$$\delta(P_i, x) \cap R_{N-j} \neq \emptyset. \quad (101)$$

The forward implication follows by restricting an accepting walk; for the converse, append arbitrary-label prefix and suffix walks of lengths i and $N - j$. Intermediate states lie in the viable layers $P_a \cap R_{N-a}$.

The pair (P_i, R_i) evolves deterministically on a finite set, so it has a common eventual threshold t and period p . Enlarge p to be divisible by every integer from 1 to s ; this does not change the threshold. Denote the stable layers by P_a^∞, R_a^∞ , for $a \in \mathbb{Z}/p\mathbb{Z}$. A total residue r is live if $P_r^\infty \cap F \neq \emptyset$. For a live r , define

$$\begin{aligned} S_{r,a} &= P_a^\infty \cap R_{r-a}^\infty, \\ \Gamma_{r,a} &= \{b \in \{0, 1\} : \Delta_b(S_{r,a}) \cap S_{r,a+1} \neq \emptyset\}. \end{aligned}$$

At sufficiently interior coordinates of a sufficiently large accepted length in residue r , these are exactly the viable state sets and coordinate alphabets. They are nonempty at every phase: an arbitrarily long accepting walk in that residue passes through every phase in its interior.

[write, batch 73] These are Part II's definitions (32), (34), (37), (38) and (39); the stable alphabet $\Gamma_{r,a}$ is Part II's $C_{r,a}$ (40), kept under the manuscript's name because C denotes a component and C_q a constant below. The enlargement of p is the one made in Part VI, Section 73.

Form a phase graph with vertices (q, a) for $q \in S_{r,a}$, and edges $(q, a) \xrightarrow{b} (q', a+1)$ whenever $(q, q') \in \Delta_b$. A word is phase-admissible at a when its successive symbols belong to $\Gamma_{r,a}, \Gamma_{r,a+1}, \dots$

Lemma 78.1 (Finite rank forces all stable words to lift; Part II, restated). *If $R(L) < \infty$, every phase-admissible word at every live residue and starting phase labels a path in the phase graph.*

[write, batch 73] This is the implication (1) $\Rightarrow$ (3) of Part II's Theorem 18.3, proved there by Theorem 18.4 for any t, p satisfying (37). The manuscript proves it again by the same pumping argument (pad a missing word to a block of length divisible by p , embed k copies between an accepted prefix and suffix, and count one internal cut per copy); that proof is not reprinted.

Lemma 78.2 (A universal strongly connected component). *In a finite phase graph with nonempty alphabets at every phase, if every phase-admissible word can be read, some strongly connected component reads every phase-admissible word at every starting phase.*

Proof. If a component misses a positive-length word v at phase b , then at any phase a prepend arbitrary admissible symbols of length congruent to $b - a \pmod{p}$. The resulting word is missing in that component at phase a . Thus failure anywhere gives an obstruction starting at every phase; if a component lacks a phase, any permitted one-letter word there supplies a positive-length obstruction.

If no component is universal, starting at any current phase concatenate one appropriately aligned obstruction for each component, updating the phase after each word. The resulting block cannot be read wholly within any component. Repeat this construction h times, where h is the number of components. Every successful walk would cross a component boundary inside each block, contradicting the fact that a walk in the condensation acyclic graph can change components at most $h - 1$ times. This contradicts universality. The universal component is cyclic, since it reads arbitrarily long words. $\square$

For a strongly connected cyclic digraph C of period d , fix cyclic classes $\kappa : C \rightarrow \mathbb{Z}/d\mathbb{Z}$, so every edge increases κ by one. We use the standard cyclic-class path property: all sufficiently large lengths a admit a walk $u \rightarrow v$ exactly when $a \equiv \kappa(v) - \kappa(u) \pmod{d}$. This follows, for example, by restricting the d th relation power to a cyclic class, where it is primitive.

Lemma 78.3 (The whole cyclic stripe is present). *Every cyclic strongly connected component of the viable phase graph has the form*

$$K = \{(v, b) : v \in C, b - a \equiv \kappa(v) - \kappa(q) \pmod{d}\} \quad (102)$$

for an original strongly connected component C , a pair $(q, a) \in K$, and the period d of C . Every internal edge between stripe pairs is present. If K is the universal component of Lemma 78.2, then $\Gamma_{r,b}$ is precisely the union of the labels on internal edges of C leaving its stripe class at phase b . Consequently the stable alphabet sequence has period dividing d .

Proof. Project paths and return paths in K to identify the original component C . A cyclic phase component projects a nonempty closed walk, so C is cyclic. Along any path the difference of phase and cyclic class stays fixed, proving one containment in (102).

For a compatible pair (v, b) , choose sufficiently long internal walks $q \rightarrow v$ and $v \rightarrow q$ whose lengths are respectively $b - a$ and $a - b$ modulo p . Such lengths exist because $d \mid p$ and the prescribed residues have the correct cyclic-class differences. Concatenate the first walk to a stable prefix reaching q at phase a , and concatenate the second before a stable suffix leaving q with complementary phase $r - a$. This gives stable prefix and suffix access to v at the required phases. Thus (v, b) is viable. The same construction applies to all intermediate compatible pairs, so these internal walks belong to the phase graph and show mutual reachability. This proves equality.

Every internal edge advances the cyclic class, hence joins stripe pairs and is present in the viable graph. Its label belongs to the stable alphabet. Conversely, universality within K reads every permitted one-letter word, so every stable symbol occurs on such an internal edge. The leaving cyclic class is determined by $b \bmod d$, giving the final periodicity statement. $\square$

This argument extracts a universal *component*, not a forward-total set of original states. No step uses deterministic images or a decrease in subset cardinality.

[write, batch 73] For a DFA, Part VI reaches a full stripe from the other side: a nonempty safe core contains a whole viable stripe with no permitted exit (Lemmas 69.1–69.3 and Theorem 69.4). Lemma 78.3 is the NFA counterpart for a universal component; the two share the core step (long internal walks of the right residue, $d \mid p$) but state different things.

79 Transferring the component bound to the whole NFA

Two quantitative facts keep a possibly large phase period from entering the rank estimate. The first uses only linear guards; the second charges entry and exit to disjoint outside states.

Lemma 79.1 (Actual slices and stable alphabets; Part VI, for NFAs). *Every accepted length $N \geq s$ has a live residue modulo p . Moreover,*

$$P_i \subseteq P_{i \bmod p}^\infty \quad (i \geq s), \quad R_j \subseteq R_{j \bmod p}^\infty \quad (j \geq s).$$

If $w \in \text{Rec}(L)_N$ and $s \leq i < N - s$, its symbol at zero-based position i belongs to $\Gamma_{N \bmod p, i \bmod p}$.

[write, batch 73] This is Part VI's Lemma 73.2, stated there for complete DFAs, with one more clause (liveness of every accepted length $N \geq s$). The manuscript's proof is Part VI's, almost word for word: a walk of at least s edges contains an elementary cycle of length $\ell \leq s$, and inserting p/ℓ copies of it adds p edges without

moving its endpoints; pumping the prefix and the suffix of a seed witness separately puts the endpoints of the witnessed letter into the stable layers. That proof uses no determinism, so it proves the lemma for NFAs as stated; applied to a whole accepting walk of length $N \geq s$ it also gives the liveness clause. It is not reprinted. The manuscript does not cite Part VI, which was written separately; only inclusion, not equality with the stable alphabets, is asserted in both.

Lemma 79.2 (Joint entry and exit budget; Part VI, for NFAs). *Fix a live residue and its universal stripe over an original component C of size c and period d . There are fixed paths from an initial state to a vertex $u \in C$ and from a vertex $v \in C$ to a final state, of lengths h, β , such that*

$$h + \beta \leq (s - c)d. \quad (103)$$

For every N in the live residue,

$$N - h - \beta \equiv \kappa(v) - \kappa(u) \pmod{d}, \quad (104)$$

and $(u, h \bmod p)$ belongs to that stripe.

[write, batch 73] This is Part VI's Lemma 73.3, with the universal stripe of Lemma 78.2 in place of Part VI's safe stripe; (103) and (104) are (93) and (94). The manuscript's proof is Part VI's, step for step: take a long accepting walk through a stripe pair, let u and v be its first and last vertices in C , shorten the parts before u and after v by deleting segments between repeated pairs (outside state, elapsed time modulo d), and note that the outside states used before entry and after exit are disjoint. That proof uses only that the stripe is a full stripe of viable pairs over an original strongly connected component, not determinism, so it proves the lemma as stated. It is not reprinted.

79.1 An imprimitive component

Suppose $d \geq 2$. Choose a cyclic class of minimum size m , so $m \leq c/d$. The internal d -step relation on this class is primitive. Its exponent e_0 obeys

$$e_0 \leq (m - 1)^2 + 1$$

by the classical primitive exponent bound; this conservative formula also holds at $m = 1$.

For any $u', v' \in C$, follow internal edges from u' for $r' < d$ steps to the chosen class, and follow internal predecessor edges backwards from v' for $t' < d$ steps to that class. There is an internal outgoing and incoming edge at every vertex, and the cyclic classes fix these two residues. Every length L of compatible residue with

$$L \geq T := d((m - 1)^2 + 1) + 2d - 2 \leq \frac{c^2}{d} + 2d - 2 \quad (105)$$

can join the two selected vertices: $L - r' - t'$ is a multiple of d and is at least de_0 . Thus every compatible pair has every compatible length at least T . The final estimate uses $(m - 1)^2 + 1 \leq m^2$.

For a hull target of length $N \geq s$, use its live residue and the paths of Lemma 79.2. Put

$$A = \max(h + T, s), \quad B = \max(\beta + T, s).$$

If $N \leq A + B$, use only singletons. Otherwise the target interval $w[A : N - B]$ follows the stable alphabet profile by Lemma 79.1. The universal stripe supplies an internal lift with the appropriate starting cyclic class. Equation (105) connects u to its first vertex in exactly $A - h$ steps and its last vertex to v in exactly $B - \beta$

steps. The stripe congruence and (104) give the two required residues. Appending the fixed outside paths produces a full accepted seed of exactly N edges matching this interval.

In both length cases the rank is at most

$$A + B + 1 \leq (s - c)d + \frac{2c^2}{d} + 4d + 2s - 3. \quad (106)$$

For fixed s, c , the function $f(d) = (s - c)d + 2c^2/d$ is convex on $[2, c]$, and

$$f(2) = c^2 + 2(s - c) \leq s^2, \quad f(c) = c(s - c + 2) \leq cs \leq s^2.$$

The second inequality uses $c \geq 2$. The maximum of a convex function on an interval is attained at an endpoint, so (106) gives

$$\rho_L(w) \leq s^2 + 6s - 3. \quad (107)$$

Unequal cyclic-class sizes and phase-dependent singleton or binary alphabets are all included. We did not recode blocks as letters or apply a binary theorem to a larger block alphabet.

79.2 A primitive component with both symbols

Now $d = 1$. Lemma 78.3 says that the stable alphabet is constant and equals the set of internal labels in C . If it is $\{0, 1\}$, the component is primitive and factor-universal, and $h + \beta \leq s - c$.

If $N < h + \beta$, singletons already suffice with at most s pieces. Otherwise put $M = N - h - \beta$ and apply Theorem 77.2 to the *actual* substring $w[h : N - \beta]$ and fixed endpoints u, v . At $M = 0$ use the empty partition. Every internally witnessed piece extends by the fixed outside paths to a complete global seed of exactly N edges. Every freely licensed singleton instead uses its witness from the global coordinate hull. The outside positions are also covered by global singleton witnesses. Thus

$$\begin{aligned} \rho_L(w) &\leq h + \beta + c^2 + 8c \\ &\leq s - c + c^2 + 8c \leq s^2 + 8s, \end{aligned}$$

since $s^2 + 8s - (s - c + c^2 + 8c) = (s - c)(s + c + 7) \geq 0$.

No extra stable-alphabet guard is required in this case: every binary symbol is in the core alphabet. We do not infer isolated-core hull membership. The permissive form of Theorem 77.2 is exactly what makes short lengths and externally licensed singletons legitimate here.

79.3 A primitive component with one symbol

The remaining alphabet is $\{\gamma\}$ for one binary symbol γ . Every internal edge then carries γ alone. Let $e_C \leq (c - 1)^2 + 1$ be the exponent of C and $M = N - h - \beta$.

If $M < e_C$, singleton fragments give

$$\rho_L(w) \leq N < h + \beta + e_C \leq c^2 + s - 3c + 2 \leq s^2.$$

The last difference is $(s - c)(s + c - 1) + 2c - 2 \geq 0$; the argument includes a negative M .

If $M \geq e_C$, choose an exact M -edge u -to- v walk. Its label is γ^M , and extending it by the outside paths gives a full seed. Put $A = \max(h, s)$ and $B = \max(\beta, s)$. For $N \leq A + B$, singletons give at most $3s - c$

pieces. Otherwise Lemma 79.1 and the unary stable alphabet show that the chosen seed matches the whole target interval $[A, N - B)$. The two singleton boundaries and this interval use at most

$$A + B + 1 \leq h + \beta + 2s + 1 \leq 3s - c + 1 \leq 3s$$

pieces.

Conclusion of the proof of Theorem 76.2. At lengths $1 \leq N < s$, use at most $s - 1$ singletons. For every nonempty hull slice of length $N \geq s$, Lemma 79.1 gives a live residue, Lemmas 78.1–79.2 give a universal stripe and the bounded exterior paths, and exactly one of the three component cases just proved applies. The empty word has rank one when accepted. The empty language has $R(L) = 1$ by definition. For $s = 1$, every nonempty slice is a coordinate product on the available loop labels and has rank one. Thus all cases satisfy the claimed $s^2 + 8s$ bound. $\square$

The qualitative eventual threshold and phase period can be large; neither appears in this estimate. The leading coefficient comes from robust primitive compression, the minimum cyclic class in the imprimitive case, and the joint outside-state budget.

80 The matching leading coefficient from a length gate

For completeness, we reprove the lower construction from the companion note (Part IV). It establishes sharpness for the unrestricted class; it does not establish sharpness within the primitive factor-universal class.

Proposition 80.1 (Part IV, restated). *For every $s \geq 4$ there is an s -state binary ε -free NFA with universal coordinate hull and*

$$R(L) = (s - 3)^2.$$

[write, batch 73] This is Part IV’s Theorem 47.1, the manuscript says so, and its proof is Part IV’s (the gate of Section 49 with two monochromatic loop states; Lemmas 49.1, 49.2 and 50.1), condensed; it is not reprinted. In it the manuscript writes H_m for Part IV’s semigroup $S_m = \langle m, m - 1 \rangle$, T_m for Part IV’s set of gate lengths H_m , and L_m for Part IV’s L_m^{gate} : the names H_m and S_m are *swapped* relative to Part IV. Section 81 refers to this gate as “the $m = s - 2$ state gate of Proposition 80.1”.

Combining Proposition 80.1 with Theorem 76.2 gives the leading coefficient in (95). The graph used for the length gate is classical; the rank argument uses the two extra monochromatic components to turn its last missing length into the number of necessary cuts. The result resolves the leading-coefficient question in the pinned report, without asserting literature-wide priority or an exact formula for $F_2(s)$. [write, batch 73] At the pin, Part II’s Question 24.2 asked only for the order of $F(s)$; the leading-coefficient question is Part IV’s Question 53.1, which the manuscript knew from Part IV’s manuscript.

81 The leading coefficient for finite alphabets

The two-cycle geometry extends beyond two symbols, but it requires a common-anchor theorem for a selected family of cycles. Pairwise intersection alone would not justify the argument. We give that structural step explicitly and keep the dependence on alphabet size visible.

Theorem 81.1 (Finite alphabet bound). *For $q \geq 2$, put $C_q = \max\{8, 2q - 4\}$. Every s -state ε -free NFA over a q -letter alphabet with finite maximal rank satisfies*

$$R(L) \leq s^2 + C_q s.$$

For every $s \geq 4$ and every $q \geq 2$ there is such an NFA with $R(L) = (s - 3)^2$, using every alphabet symbol. Hence

$$(s - 3)^2 \leq F_q(s) \leq s^2 + \max\{8, 2q - 4\}s, \quad \lim_{s \rightarrow \infty} \frac{F_q(s)}{s^2} = 1$$

for each fixed q . The same ratio tends to one along any integer sequence $q = q(s) \geq 2$ with $q(s) = o(s)$. In particular the explicit upper bound $s^2 + 8s$ holds for $2 \leq q \leq 6$.

81.1 A common vertex for a small cycle family

Work in an n -vertex digraph whose only elementary cycle lengths are coprime integers $j < k$, where $j > n/2$ and $j \geq 3$. Choose $t < j$ elementary cycles, allowing repetitions, of which a have length j and b length k .

Lemma 81.2 (Uniqueness of the cycle counts). *Every decomposition of the sum of these cycles' integral edge-incidence vectors into elementary cycles has exactly a cycles of length j and b of length k .*

Proof. Any such decomposition gives $aj + bk = Aj + Bk$ with $A, B \geq 0$. Coprimality implies $A = a + hk$ and $B = b - hj$ for an integer h . If $h > 0$, then $B < 0$ because $b < j$; if $h < 0$, then $A < 0$ because $a < j < k$. Therefore $h = 0$. Every nonnegative integral circulation decomposes into elementary cycles by successively removing a cycle from its positive support. $\square$

Lemma 81.3 (Small family common vertex). *The t selected cycles have a common vertex.*

Proof. Sum their edge-incidence vectors as an Eulerian directed multigraph. At vertex v the indegree and outdegree both equal d_v , the number of selected cycles containing v . If there were no common vertex, then $\Delta := \max_v d_v \leq t - 1$. Add $\Delta - d_v$ new loops at each vertex, making all indegrees and outdegrees Δ .

Replace every vertex by a left and a right copy and every directed edge by the corresponding bipartite edge. The resulting bipartite multigraph is Δ -regular. For any left vertex subset, counting its incident edges proves Hall's condition, so a perfect matching exists; deleting it and repeating decomposes the multigraph into Δ perfect matchings. Each matching is a directed cycle cover. Its added loops are isolated one-vertex components. Every other cycle uses original edges and has at least $j > n/2$ vertices, so each cover contains at most one such cycle. Removing the added loops therefore decomposes the original circulation into at most $\Delta < t$ elementary cycles, contradicting Lemma 81.2. $\square$

Fix a common vertex z . Let U_j, U_k be the edge unions of the selected cycles of their respective lengths, omitting an empty group. Each nonempty union is strongly connected, since its constituent cycles contain z . Every elementary cycle of U_j has length j : if a k -cycle existed, subtract one copy from the sum of the selected j -cycle circulations and decompose the nonnegative remainder, contradicting Lemma 81.2. Similarly U_k has only length k . Thus their periods are j and k , respectively.

Index the cyclic classes of each group from z . Every selected cycle visits each of its group's classes exactly once. Write $m_j = |V(U_j)|$, $m_k = |V(U_k)|$, and $m = |V(U_j) \cup V(U_k)| \leq n$. Let S_j, S_k be the vertices constituting singleton classes in the corresponding groups. Counting non-singleton classes gives

$$|S_j| \geq 2j - m_j, \quad |S_k| \geq 2k - m_k. \quad (108)$$

A vertex is on every selected cycle of one length if and only if it is a singleton-class vertex of that group's union: every cycle meets the class once, and every vertex of the union belongs to a selected cycle.

81.2 Aligning the two class systems

Assume both length groups are nonempty. For any vertex v common to U_j, U_k , its class indices $p_j(v) \in \{0, \dots, j-1\}$ and $p_k(v) \in \{0, \dots, k-1\}$ are its exact positions after z on every selected cycle of the corresponding group that contains it. We claim

$$p_k(v) - p_j(v) \in \{0, k-j\}. \quad (109)$$

For $v \neq z$, choose one cycle of each length containing v and cross-splice their z -to- v prefixes and v -to- z suffixes. The two resulting nonempty closed walks have lengths $j + p_k(v) - p_j(v)$ and $k - p_k(v) + p_j(v)$. Their total is $j + k$. Decomposition and Lemma 81.2 give exactly one j -cycle and one k -cycle in total. Each nonempty walk must supply one, proving (109). For $v = z$, both indices are zero and the claim is immediate.

Lemma 81.4 (Common anchor bound). *If both length groups occur, the selected cycles have at least $2j + k - 2n$ common vertices.*

Proof. Their common vertices are exactly $S_j \cap S_k$. Consider $v \in S_j \cap V(U_k)$ not in S_k , and another vertex w in its k -class. We first prove $w \notin V(U_j)$.

Otherwise (109) says that $p_j(v), p_j(w)$ either agree or differ in absolute value by $\delta = k - j$. Equality is impossible because the j -class of v is a singleton. Also $0 < \delta < j$, since $k \leq n < 2j$. A selected j -cycle containing w also contains v , and in the forward order of their class indices supplies a path of length δ between them. The group U_k supplies a return path of exactly k edges: a shortest path between the distinct same-class vertices has positive length divisible by k and at most $m_k - 1 \leq n - 1 < 2k$.

Concatenation gives a closed walk of length $k + \delta = 2k - j$. This length is not a nonnegative combination of j, k . It can contain at most one k -term. With no k -term, divisibility would force $j \mid 2k$, hence $j \leq 2$; with one, it would force $j \mid k$. Both contradict the hypotheses. Decomposing the closed walk into elementary cycles is therefore impossible, proving the claim.

Two vertices of S_j cannot be in the same k -class, since the claim applied to one would exclude the other from U_j . Each vertex of $S_j \cap U_k$ not in S_k can consequently be charged injectively to a vertex of $U_k \setminus U_j$ in its distinct k -class. The vertices of S_j absent from U_k are at most $|U_j \setminus U_k|$. Using (108),

$$\begin{aligned} |S_j \cap S_k| &\geq |S_j| - |U_k \setminus U_j| - |U_j \setminus U_k| \\ &\geq (2j - m_j) - (m - m_j) - (m - m_k) \\ &= 2j + m_k - 2m \geq 2j + k - 2n. \end{aligned}$$

□

81.3 Compression with one branch for every symbol

The exponent reduction preceding (99) does not depend on alphabet size. Thus for a primitive factor-universal q -letter graph it leaves only the hard case $n \geq 6$, $n/2 < j < k \leq n$, $\gcd(j, k) = 1$, and $e \leq jk$. Choose one monochromatic elementary cycle for each symbol. If $n \geq 2q$, then $q < j$ and the preceding lemmas apply to all selected cycles at once.

First suppose all chosen cycles have length $\ell \in \{j, k\}$; let r be the other length. Their union U is strongly connected and has only elementary cycle length ℓ , by Lemmas 81.3 and 81.2. It has ℓ cyclic classes, of which at least $b \geq 2\ell - m \geq 1$ are singletons, where $m = |V(U)| \leq n$. These singleton vertices lie on every selected cycle. Between consecutive such anchors, each symbol has a monochromatic path of the same length, equal to the cyclic-class distance. With only one anchor, the region is the full positive-length return traversal of length ℓ .

In each equal-length region give its first position charge zero and its other positions charge one. A fragment may cross an anchor without a cut. Cutting only at symbol changes internal to a region allows each fragment to choose the relevant monochromatic branch independently, replacing whole regions in its complete witness even when the fragment begins or ends inside one. Hence charge plus one pays for the fragments in a segment. This is the same fixed-endpoint argument as Lemma 77.5; disjoint interiors of the branches are unnecessary for valid walks.

A lap has charge $h = \ell - b \leq m - \ell$. Set $\mu = h/\ell < 1$. In a repeated selected ℓ -cycle, a segment of length x has charge at most $\mu x + h$: remove full periods beginning at the segment's own phase and bound its remaining proper-period charge by a complete lap. Choose an elementary r -cycle meeting a selected ℓ -cycle at z' , and choose a singleton anchor y . Fixed paths $u \rightarrow z'$, $z' \rightarrow y$ along that cycle, and $y \rightarrow v$ have combined length $D \leq 2n + \ell - 3$. Above the Frobenius threshold write

$$M - D = \alpha r + \beta \ell, \quad 0 \leq \alpha < \ell, \quad \beta \geq 0,$$

and use the complete walk formed by these connectors, the single r -cycle block at z' , and the repeated ℓ -cycle block at y . If the threshold fails, $M < jk + 2n - r - 2 < n^2 + 2n$, so singletons suffice.

Cover the entire target if $M \leq 2e$, or the two e -position boundary windows and a universal middle fragment otherwise. Their combined length is $T \leq 2e$. Their total intersection K with the single r -block is less than jk . The repeated ℓ -block meets at most two window segments, so the number of pieces is at most

$$\begin{aligned} \mu T + (1 - \mu)K + D + 2h + 3 &\leq jk(1 + \mu) + D + 2h + 3 \\ &= r(\ell + h) + D + 2h + 3 \\ &\leq rm + D + 2h + 3 \leq n^2 + 4n. \end{aligned}$$

For the last estimate use $r, m \leq n$, $D \leq 2n + \ell - 3$, and $h \leq n - \ell$.

Now suppose both cycle lengths occur. Let b be the number of vertices common to all selected cycles. Lemma 81.4 gives $b \geq 2j + k - 2n$. If $b \leq 2$, then $2j + k \leq 2n + 2$ and the elementary bound suffices:

$$2e + 1 \leq 2jk + 1 \leq 4j(n + 1 - j) + 1 \leq (n + 1)^2 + 1 \leq n^2 + 8n.$$

If $b \geq 3$, the pairwise order lemma puts all anchors in the same cyclic order on every selected cycle. Within a length group their class indices fix every anchor-to-anchor path length. Comparing one j -cycle with one k -cycle and coarsening their pairwise bubbles to these anchors shows that exactly one region has unequal lengths, differing by $k - j$. All other regions offer an equal-length monochromatic branch for every symbol. This asserts only a property of the selected paths, not a global bubble decomposition of the whole graph.

Charge the equal regions as before. Charge every position of the unequal region one, with an additional unit on its first position. This pays for its singleton fragments and the necessary interfaces. A segment again needs at most its charge plus one pieces. The two lap charges are

$$c_j = j - b + 2, \quad c_k = k - b + 2, \quad \mu = \frac{k - b + 2}{k} \leq 1.$$

Their mean charges are at most μ . A segment of length x in either repeated cycle has charge at most $\mu x + n$, since its residual proper-period charge is at most one lap, which is at most n .

Choose an anchor y and fixed entry and exit paths of total length $D \leq 2n - 2$. Above the Frobenius threshold, use $M - D = \alpha j + \beta k$ and the complete ambient walk with a j -block followed by a k -block at y . Below the threshold, singletons give fewer than $n^2 + 2n$ pieces. The two boundary windows meet these two consecutive repeated-cycle blocks in at most three nonempty periodic segments. That count is used for the residual-charge overhead $3n$. The junction of the blocks is an anchor, so it adds no fragment cut: the combined repeated-cycle block meets the windows in at most two pieces of charge-plus-one bookkeeping. Including connectors and the possible middle fragment gives

$$\begin{aligned} 2\mu e + 3n + D + 3 &\leq 2j(k - b + 2) + 5n + 1 \\ &\leq 4j(n - j + 1) + 5n + 1 \\ &\leq n^2 + 7n + 2 \leq n^2 + 8n. \end{aligned}$$

This proves the robust primitive bound $n^2 + 8n$ for $n \geq 2q$.

For $n < 2q$, use the general exponent bound $2e + 1 \leq 2(n - 1)^2 + 3$. For $q \geq 3$ and $2 \leq n \leq 2q - 1$,

$$n^2 + (2q - 4)n - (2(n - 1)^2 + 3) = n(2q - n) - 5 \geq 0.$$

The concave quadratic is minimized at an interval endpoint, where it is at least $2q - 6 \geq 0$. The one-state case has rank one, and $q = 2$ is covered by Theorem 77.2. We have therefore proved the robust primitive bound $n^2 + C_q n$ for every n .

81.4 Transfer and the lower construction

The proofs of the stable-word, universal-component, stripe, guard, entry–exit, and imprimitive-component results in Sections 78–79 work over any finite alphabet: replace $\{0, 1\}$ in the stable-alphabet definition by Σ . The imprimitive estimate $s^2 + 6s - 3$ is within the desired bound because $C_q \geq 8$.

For a primitive universal component of size c , let B be its nonempty internal alphabet and let $h + \beta \leq s - c$ be the fixed exterior-path budget. Over B , the robust bound is at most $c^2 + C_q c$, since $C_{|B|} \leq C_q$ when $|B| \geq 2$. For $|B| = 1$ the same estimate holds directly: a target has the unique unary label, and an exact internal witness exists for lengths at least the exponent, while shorter lengths cost at most $(c - 1)^2 + 1$ singletons.

If $c = s$, all seed symbols lie in B , so apply that robust bound directly. If $c < s$, put $A = \max(h, s)$ and $Z = \max(\beta, s)$. When $N \leq A + Z$, use singletons. Otherwise project $w[h : N - \beta]$ to B by retaining its symbols in B and replacing the others by a fixed symbol in B . Apply the robust primitive theorem with the required fixed endpoints and full internal length. Restrict the resulting partition to the contiguous global interval $[A, N - Z]$; its piece count does not increase. By Lemma 79.1, the original target already lies in B there. Thus all retained internal witnesses match the actual target and extend by the fixed exterior paths to accepted length- N seeds. Licensed and outside singletons use their global hull witnesses. The total cost is at most

$$h + \beta + 2s + c^2 + C_q c \leq s - c + 2s + c^2 + C_q c \leq s^2 + C_q s.$$

Indeed the last slack is $(s - c)(s + c + C_q - 1) - 2s \geq C_q - 2 \geq 0$ for $c \leq s - 1$. Short lengths, the empty word, and the empty hull are handled as in the binary proof. This establishes the upper bound of Theorem 81.1.

For the lower bound, keep the $m = s - 2$ state gate of Proposition 80.1 and label every gate edge by all symbols. Partition $\Sigma = A \sqcup B$ into two nonempty classes. Give one extra isolated state all loop labels in A and the other all loop labels in B , with both states initial and final as before. At a gate length every word is accepted. At a missing gate length N , the seeds are exactly $A^N \cup B^N$, so the hull is Σ^N and an interval is admissible exactly when its symbols belong to one class. Rank is therefore the number of class runs.

Alternating one symbol from each class at the last exceptional length $(m - 1)^2$ proves exact rank $(s - 3)^2$. The quantitative sandwich implies both asymptotic statements.

No alphabet-independent additive constant is established here. When q is comparable to s , the displayed upper bound alone does not give a leading coefficient of one; for unrestricted alphabet growth that question remains open in this article.

82 A sharp linear family and its slice maxima

The quadratic length-gate family is asymptotically extremal, but is not the best lower bound for small state counts. A different construction attains the general exponent estimate $2e + 1$ with a linear exponent, and supplies the lower witnesses for the three- and four-state extrema below.

Proposition 82.1 (Shortest access and exit gates). *For every $s \geq 3$ there is a primitive factor-universal binary s -state NFA with sole initial state u , sole final state v , exponent $e = s - 1$, and*

$$R(L) = 2s - 1 = 2e + 1.$$

More precisely, for every $N \geq 1$ its slice maximum is

$$M_L(N) := \max\{\rho_L(w) : w \in \text{Rec}(L)_N\} = \min\{N, 2s - 1\},$$

and 0^N attains that maximum. Its empty word is not accepted.

Proof. Use distinct states u, v, b , and put a 1-loop at every state. The only 0-edges are $u \rightarrow v$ and $b \rightarrow b$. For $s = 3$, add 1-edges $v \rightarrow b$ and $b \rightarrow u$. For $s \geq 4$, introduce states $x_1, \dots, x_{s-3}$ and add precisely these 1-edges:

$$v \rightarrow x_1, \quad b \rightarrow x_1, \quad x_i \rightarrow x_{i+1} \quad (1 \leq i < s - 3), \quad x_{s-3} \rightarrow u, \quad x_{s-3} \rightarrow b.$$

There are no further edges. Take $I = \{u\}$ and $F = \{v\}$.

The underlying graph is strongly connected. Every walk from u to b must use $u \rightarrow v$ and then the full corridor; every walk from b to v must traverse the full corridor to u and then use $u \rightarrow v$. Thus

$$\text{dist}(u, b) = \text{dist}(b, v) = s - 1. \tag{110}$$

For $s = 3$ this is immediate from its directed three-cycle. Every ordered pair is connected by a simple path of at most $s - 1$ edges. The loops pad it to exactly $s - 1$ edges, and (110) rules out a smaller exponent. Hence $e = s - 1$. The two loops at b establish factor universality, so the central-fragment argument gives $R(L) \leq 2s - 1$.

At each zero-based coordinate i of a positive length N , the seed $1^i 0 1^{N-i-1}$ is accepted: wait at u , use $u \rightarrow v$, and wait at v . Consequently 0^N belongs to the hull. Any fragment of the target 0^N of length at least two must use only the 0-loop at b , since the other 0-edge ends at v and no 0-edge leaves v . By (110), such a fragment can begin only at a boundary at least $s - 1$ and end only at a boundary at most $N - s + 1$.

If $N < 2s - 1$, no nonsingleton zero fragment fits, so $\rho_L(0^N) = N$. If $N \geq 2s - 1$, the first and last $s - 1$ positions each force singleton pieces, and the remaining nonempty middle requires at least one piece. Thus $\rho_L(0^N) \geq 2s - 1$. The singleton bound and the established global upper bound give equality in both cases, proving the slice formula. $\square$

This also identifies the exact maximum within the useful universal-loop class. If an s -state NFA has a state carrying every symbol loop and lying on an accepting walk, choose shortest access and exit paths of lengths $a, b' \leq s - 1$. At length $N > a + b'$, the target's middle interval $[a, N - b')$ is read at that state; its witness is completed by the two fixed paths. The outside positions are hull singletons. Shorter targets use only singletons. Hence $R(L) \leq a + b' + 1 \leq 2s - 1$. Proposition 82.1 attains this for each $s \geq 3$. This observation explains a limitation of searches that always insert a universal-loop state: they cannot discover superlinear rank in the full primitive class.

For binary NFAs the two ordinary lower constructions together give

$$F_2(s) \geq \max\{2s - 1, (s - 3)^2\} \quad (s \geq 4).$$

Neither construction alone is asserted to give the exact general extremum.

83 An exact algorithm for the maximal rank

The asymptotic state bound and the exact value for one automaton are different questions. We now give an effective exact algorithm for any fixed finite alphabet, together with a complete finite witness-length cutoff. No guessed experimental horizon is used.

83.1 The greedy parser is optimal

[write, batch 73] Lemma 83.1, the parser after it and the periodic-middle graph below are Part V's Lemma 57.1, its parser of Section 58 and its reset graph G_{r,a_0} of Section 60, which the manuscript cites elsewhere but not here. The lemma is printed as a restatement with the manuscript's proof as a second proof; the parser is the same procedure. What is new here is the treatment of the two boundaries, A_r and D_r , the positive-cycle test and the cutoff.

Fix a nonempty hull slice of length N and put $S_i = P_i \cap R_{N-i}$. An interval is admissible exactly when its label can be read from S_i to S_j through the intervening viable layers; this is equivalent to (101). In particular, restriction of an admissible interval is admissible, and every singleton of a hull word is admissible.

Lemma 83.1 (Greedy optimality; Part V, restated). *Partition a nonempty hull word by repeatedly taking its longest admissible prefix from the current boundary. This partition has minimum size.*

Second proof (manuscript 21). Let g_k be the greedy boundaries and let q_k be those of any comparison partition. Inductively $g_k \geq q_k$ as long as the greedy parser has not already finished. If $g_{k-1} \geq q_k$, the claim is immediate. Otherwise the induction hypothesis gives $q_{k-1} \leq g_{k-1} < q_k$, so $[g_{k-1}, q_k)$ is a restriction of the comparison partition's k th admissible interval. The longest greedy extension reaches at least q_k . Thus the greedy parser finishes in no more pieces than any comparison partition. $\square$

The parser stores the nonempty subset Z of possible endpoints of its current fragment. Initially $Z = S_0$ and its rank counter is 1. At zero-based coordinate i , on letter a , compute

$$Z' = \Delta_a(Z) \cap S_{i+1}.$$

If this is nonempty, replace Z by Z' . Otherwise cut immediately before that letter, add one to the rank counter, and restart with

$$Z = \Delta_a(S_i) \cap S_{i+1},$$

which is nonempty by coordinate admissibility. The subset is exactly the possible endpoints of the current greedy fragment. A failed extension cannot be rescued by later letters, so this procedure implements Lemma 83.1. A nonempty target therefore has rank 1 plus its reset count. Taking maxima of this deterministic parser output does not interchange a minimum and a maximum.

At a single length N , dynamic programming keeps for each current nonempty subset the largest rank counter attained there and advances over all coordinate-permitted letters. This computes $M_L(N)$ without enumerating target words; a maximizing target is recovered by retaining predecessor choices.

83.2 Finite boundary operators and the periodic middle

Choose any common eventual transient $t \geq 0$ and period $p \geq 1$ for (P_i, R_i) . In this section it is unnecessary to enlarge p to a multiple of all cycle lengths. They can be found by following the deterministic orbit of pairs of subsets until its first repetition. Stable phase values are extended periodically to all integer phase arguments.

For a live total residue r , every $N \geq 2t$ with $N \equiv r \pmod{p}$ has the following viable layers, independent of the particular N :

$$\begin{aligned} \text{left boundary:} & \quad P_i \cap R_{r-i}^\infty \quad (0 \leq i \leq t), \\ \text{middle phase:} & \quad S_{r,a} = P_a^\infty \cap R_{r-a}^\infty, \\ \text{right boundary:} & \quad P_{r-t+j}^\infty \cap R_{t-j} \quad (0 \leq j \leq t). \end{aligned}$$

The descriptions agree at the joins. A coordinate alphabet consists exactly of letters with an edge between its two consecutive viable layers. All such alphabets are nonempty on a live slice.

Run the parser on all coordinate-permitted t -letter prefixes, retaining for each terminal subset Z the largest rank-so-far $A_r(Z)$. This subset is at phase $t \bmod p$. When $t = 0$, the initial value 1 is assigned to the initial viable subset. Similarly, for each nonempty subset at phase $r - t \bmod p$, let $D_r(Z)$ be the largest additional reset count in the final t coordinates. These finite operators can be computed by the same subset dynamic programming. Every eligible nonempty Z can process the suffix, using resets if needed, so each $D_r(Z)$ is finite. Ineligible or unattainable values are $-\infty$.

The periodic-middle graph has vertices (a, Z) , where $a \in \mathbb{Z}/p\mathbb{Z}$ and $\emptyset \neq Z \subseteq S_{r,a}$. For each $c \in \Gamma_{r,a}$, advance to phase $a + 1$ and put a directed edge of weight

$$\begin{cases} 0, & \text{with target } \Delta_c(Z) \cap S_{r,a+1}, \quad \text{if that set is nonempty,} \\ 1, & \text{with target } \Delta_c(S_{r,a}) \cap S_{r,a+1}, \quad \text{otherwise.} \end{cases}$$

Initialize the phase- t vertices with their A_r values and discard unreachable vertices. [write, batch 73] With $t \bmod p$ as starting phase this is Part V's graph $G_{r,t \bmod p}$, whose vertex bound $p(2^s - 1)$ is the same. Every retained vertex has an outgoing edge, and the graph has at most

$$V = p(2^s - 1) \tag{111}$$

vertices.

A complete word of length $N \geq 2t$ in residue r is exactly a permitted prefix, a middle path of length $N - 2t$ ending at phase $r - t$, and a permitted suffix. Its rank is the prefix value plus the path weight plus its suffix resets. Conversely every such combination gives a hull word, because a coordinate hull independently permits each of its coordinate letters. The terminal phase enforces the required congruence on the total length; there is no additional hidden congruence restriction.

Theorem 83.2 (Exact graph algorithm). *The maximal rank is unbounded if and only if one of the reachable periodic-middle graphs contains a directed cycle of positive total weight. Otherwise, contracting its strongly*

connected components and optimizing on the resulting directed acyclic graph computes the exact maximal rank over all sufficiently long slices. Together with the slices of lengths less than $2t$ and the empty-word convention, this computes $R(L)$ exactly.

Proof. Every reachable middle vertex can reach a vertex at phase $r - t$ by following arbitrary outgoing edges for the appropriate number of steps, chosen below p . The suffix can then be processed. Thus every reachable directed cycle is co-accessible to a complete word. A positive-weight cycle may be repeated between fixed access and exit paths. Its length is a multiple of p , so these complete words retain the same residue and boundary operators; their ranks grow without bound.

Conversely, if no reachable positive-weight cycle exists, every edge within a reachable strongly connected component has weight zero: a positive edge there could be completed to a positive cycle, since all weights are nonnegative. Contract these components. At a contracted vertex retain the maximum of the initial values A_r and, separately, the maximum eligible terminal value D_r among its members. Arbitrary entry and exit vertices in one component can be connected at zero cost. These maxima are therefore compatible with any incoming and outgoing component edges, even when attained at different original vertices. A longest-path dynamic program in the condensation graph gives exactly the largest sum of initial, transition, and terminal values. Its paths correspond to finite original paths, and hence to complete hull words as above. The finitely many short slices and the accepted empty word finish the calculation. If the hull is empty our definition gives $R(L) = 1$. $\square$

Theorem 83.3 (Complete finite witness cutoff). *If $R(L) < \infty$, every hull word of length $N \geq 2t$ has the same rank as a hull word in the same total residue of length at most*

$$2t + p(2^s - 1) - 1.$$

Consequently, if the hull is nonempty, its maximal rank is attained at some length at most

$$H = \max\{1, 2t + p(2^s - 1) - 1\}, \tag{112}$$

with length zero included when the empty word is accepted.

Proof. Take the target's middle parser path. By Theorem 83.2, all its directed cycles have weight zero. Repeatedly delete closed portions between repeated vertices. The resulting simple path has at most $V - 1$ edges and the same endpoints and total weight. Each removed length is divisible by p , so the terminal and total phases are unchanged. Keep the original t -letter prefix and suffix and the remaining middle edge labels. At the new total length all boundary layers are unchanged and each retained middle letter is still coordinate-permitted. This reconstructs a hull word with the same prefix parser value, middle reset count, final subset, and suffix reset count. Its rank is unchanged and its length is at most $2t + V - 1$.

Shorter words already satisfy the bound, since $V \geq 1$. The artificial lower bound $H \geq 1$ only avoids a zero positive-length horizon in trivial cases. For an empty hull, the value $R(L) = 1$ is a convention and is not asserted to have a witness. $\square$

[write, batch 73] This complements Part II's Corollary 19.4, which bounds the length of high-rank witnesses when stabilization fails, and bears on Part II's Question 24.4; neither is the other.

An alternative exact implementation first decides boundedness using stable-word universality and then computes all slice maxima up to H . Universality is decidable by the finite phase/subset graph that starts from each entire viable layer and checks whether a permitted word can reach the empty set.

Corollary 83.4 (A transient bound for finite rank; Part II, restated). *For any common eventual transient t of (P_i, R_i) ,*

$$R(L) < \infty \implies R(L) \leq 2t + 1.$$

[write, batch 73] This is the bound (42) of Part II’s Theorem 18.3, which holds for any t, p satisfying (37); Part V uses it the same way (Corollary 64.2). The manuscript’s proof is Part II’s (one central fragment and $2t$ singleton boundaries) and is not reprinted.

This equivalence also proves the boundedness test just described: failure of stable-word universality yields unbounded ranks by the obstruction repetition, and success yields the transient bound. Although t is useful for exact computation and pruning, it need not be a sharp state-count parameter; the compression theorem avoids paying for a potentially large transient twice.

84 Computer assisted binary extrema through four states

Here $F_2(s)$ is the binary extremal function. Its argument counts states; it is not the alphabet index. All upper census conclusions here concern exactly two symbols. The exact algorithms of Section 83 apply to other fixed finite alphabets, but these binary censuses do not determine their extrema.

Theorem 84.1 (Exact small binary extrema). *For binary ε -free NFAs with at most s states,*

$$F_2(1) = 1, \quad F_2(2) = 2, \quad F_2(3) = 5, \quad F_2(4) = 7.$$

The upper bounds for two and three states are established by exhaustive finite computation with independent implementations. The four-state upper bound combines Corollary 83.4 with an exhaustive pruned computation. The lower bounds have explicit witnesses.

84.1 One to three states

For one state every nonempty accepted slice is the full product of its loop alphabet, so the maximal rank is one. For two states, use initial state v and final state u , 1-loops at both states, and 0-edges $v \rightarrow v$ and $v \rightarrow u$, with no other edges. This accepts exactly the binary words containing at least one 0: switch to u at the last zero. A seed has rank one. At a length $N \geq 2$, the only nonseed hull target is 1^N , and it has rank two: split it into two nonempty intervals and place a zero outside each interval in its witnessing seed. It is not itself a seed. At length one, only 0 belongs to the hull. Thus this NFA has $R(L) = 2$, attained by 11.

For three states Proposition 82.1 already gives rank five with the target 0^5 . The census supplies the upper bound.

There are $2^{2s^2}(2^s - 1)^2$ labeled binary transition pairs with nonempty initial and final sets. Presentations with an empty endpoint set recognize the empty language and cannot raise $F_2(s)$. The complete two- and three-state distributions are shown below. Here “empty” denotes an empty hull, recorded as raw maximum 0 by the programs; under this article’s floor convention it has $R(L) = 1$. Every numbered rank column excludes

those empty presentations.

outcome	$s = 2$	$s = 3$
empty	128	270,336
$R = 1$	1,612	8,680,660
$R = 2$	188	2,239,104
$R = 3$	0	111,984
$R = 4$	0	2,016
$R = 5$	0	72
$R = \infty$	376	1,540,884
total	2,304	12,845,056

One implementation quotients transition pairs by state permutations and exchange of the binary letters, enumerates all nonempty endpoint sets, and weights each result by its exact transition-pair orbit size. For $s = 3$ there are 22,284 transition-pair representatives. It uses the complete cutoff and slice dynamic programming. A second implementation uses full Boolean-matrix eventual periodicity, exact boundary operators, positive-cycle detection, and periodic-middle path optimization. A third, independently written implementation directly enumerates every labeled transition pair and endpoint choice, using the paired-layer orbit and explicit boundary/core graph with strongly connected components. All three agree pointwise on all 12,847,360 two- and three-state presentations.

The finite cutoff is at most 6 in the finite two-state cases and at most 20 in the finite three-state cases. An independent literal oracle enumerates accepted seeds and their actual aligned substrings, then minimizes over all interval partitions. It corroborates the parser on 208,712 two-state hull targets and 147,264 targets from all 72 three-state rank-five maximizers. These cross-checks are additional evidence, not a replacement for the complete enumeration.

84.2 A complete four-state pruning rule

A brute-force full rank distribution is unnecessary for the four-state maximum. For each underlying relation E and nonempty I, F , let t be the least transient of the paired orbit

$$(P_0, R_0) = (I, F), \quad (P_{i+1}, R_{i+1}) = (E(P_i), E^\top(R_i)).$$

This orbit depends only on E, I, F , before the present edges are labeled. By Corollary 83.4, any finite-rank counterexample to $R(L) \leq 7$ must have $t \geq 4$. All labelings of every discarded pair with $t \leq 3$ are therefore covered by an ordinary proof.

Enumerate the 2^{16} four-state underlying relations, including loops, up to state permutation. There are 3,044 graph representatives. For each one, examine all 15^2 nonempty ordered endpoint-set pairs and retain exactly those with $t \geq 4$. This leaves 411 graphs and 23,228 triples (E, I, F) .

For each retained graph with m present edges, enumerate its 3^m assignments of each edge as 0-only, 1-only, or both. These are exactly the binary transition pairs with union E . Retain one from each exchange of the two letters, using their integer matrix encodings to impose $\Delta_0 \leq \Delta_1$. Exactly one labeling is fixed by that exchange, namely the all-doubly-labeled graph. Thus each endpoint pair contributes $(3^m + 1)/2$ cases. The total is independently predicted as

$$\frac{91,368,216 + 23,228}{2} = 45,695,722.$$

Every possible counterexample is covered: move its underlying graph to the chosen state-permutation representative and transport its labels and endpoint sets by the same permutation. All those endpoint sets

and labelings are examined; exchange of the letter names preserves rank. It is unnecessary to quotient endpoint sets by graph stabilizers, since duplicate isomorphic cases do not affect a maximum. No trimming, connectedness, primitivity, or singleton-endpoint assumption is imposed.

The complete retained scan has the following outcome. These are unweighted counts inside the pruned graph-representative and letter-exchange search, *not* a distribution over all labeled four-state NFAs.

outcome	count	outcome	count
empty	186,704	$R = 1$	14,622,000
$R = 2$	15,145,667	$R = 3$	3,127,160
$R = 4$	202,894	$R = 5$	3,837
$R = 6$	80	$R = \infty$	12,407,380

The counts sum to 45,695,722 and the largest finite rank in the retained class is 6. In particular there is no retained counterexample to rank at most 7. The discarded cases already satisfy that bound, and Proposition 82.1 attains 7 with four states. Thus $F_2(4) = 7$. Padding with isolated states makes the conclusion valid for the at-most-state convention.

An independent candidate generator uses orbit marking and Floyd cycle detection, rather than the producer's minimum-under-permutations test and indexed orbit table. It regenerates the same 23,228-record candidate list. Its independently written label enumeration and boundary/core graph rank engine completes all 45,695,722 retained cases and reproduces every displayed count. Finite inputs are checked with the complete graph optimization; unboundedness is tested by the proved interior-universality criterion, with 12,492 deterministic unbounded samples additionally checked by the full graph. All paired orbits have $t \leq 10$, $p \leq 4$, and cutoff expression $2t + 15p - 1 \leq 59$; retained graphs have at most ten edges. These bounds also validate the finite arrays and integer encodings used by the implementations.

The exact small-state conclusions are computer-assisted mathematical results with compact receipts and independently reproducible source. They are not proof-assistant certificates, and they do not claim $F_2(s)$ for any $s \geq 5$.

85 Sharper bounds from common paths

Theorem 85.1. *Let L be recognized by an n -state, ε -free, factor-universal binary NFA whose underlying digraph is primitive and has a Hamilton cycle. For arbitrary nonempty initial and final sets,*

$$R(L) \leq n^2 + n + 1.$$

The common-path argument gives a smaller additive term than Theorem 77.2 in this restricted class.

85.1 The central fragment and boundary compression

For an underlying digraph with edge relation E , let

$$\begin{aligned} P_t &= \{q : \text{an } E\text{-walk of length } t \text{ goes from } I \text{ to } q\}, \\ R_t &= \{q : \text{an } E\text{-walk of length } t \text{ goes from } q \text{ to } F\}. \end{aligned}$$

These are exact-length layers, with labels disregarded.

Lemma 85.2. *For a primitive factor-universal NFA of exponent e with nonempty initial and final sets, one has $P_e = R_e = Q$. Every hull word w of length $N > 2e$ has an admissible central interval $[e, N - e]$. Consequently $R(L) \leq 2e + 1$.*

Proof. Since I, F are nonempty, the definition of exponent gives both layer equalities. Factor universality realizes $w[e : N - e]$ as a labeled walk. Prepend a length- e walk from I to its first vertex and append a length- e walk from its last vertex to F . Choose any available edge labels on the added walks. The result is a complete accepted seed matching the central interval. Singleton hull witnesses cover the two boundaries. If $1 \leq N \leq 2e$, the singleton bound suffices; the empty word has rank at most one when present. $\square$

The next lemma is the quantitative step. A directed path is simple, so its vertices are distinct.

Lemma 85.3 (Common path edge charge). *Suppose a labeled digraph on n vertices has a directed path H of $c \geq 1$ edges with these properties:*

1. *Every edge of H carries both letters.*
2. *No edge outside H enters or leaves an internal vertex of H .*
3. *Give weight one to the first edge of H and to each edge outside H , and weight zero to the other edges of H . Every elementary directed cycle has mean weight at most μ , where $0 \leq \mu \leq 1$.*

Fix a hull target of length N and any underlying accepting walk of length N . Every segment of this walk of length $\ell \geq 1$ yields a partition of the corresponding target interval into at most

$$\mu\ell + (1 - \mu)(n - 1) + 1 \quad (113)$$

admissible pieces. An empty segment needs no pieces.

Proof. Delete elementary cycles successively from the segment, leaving a simple path of $m \leq n - 1$ edges. Its weight is at most m . The removed cycles have total weight at most $\mu(\ell - m)$, so the original segment's weight W satisfies

$$W \leq \mu\ell + (1 - \mu)m \leq \mu\ell + (1 - \mu)(n - 1). \quad (114)$$

This cycle deletion is used only to count weight; fragments remain on the original accepting walk.

Make each position on a non- H edge a singleton, using its hull witness, and each nonempty maximal run of H edges one piece. Each run is admissible: on that run choose the target's letters on the doubly labeled edges, and choose any labels on the other edges of the same complete accepting walk. The endpoints and total length are unchanged.

Except for a possible initial truncated run, every run enters through the first edge of H , because internal vertices admit no outside entry. Charge each such run to that edge. Charge each non- H singleton to its own edge. The number of pieces is at most $W + 1$, proving (113). This reasoning also covers $c = 1$. $\square$

Corollary 85.4. *Under the hypotheses of Lemmas 85.2 and 85.3,*

$$R(L) \leq 2\mu e + 2(1 - \mu)(n - 1) + 3 = 2\mu(e - n + 1) + 2n + 1. \quad (115)$$

Proof. If $N > 2e$, apply Lemma 85.3 separately to the first and last e edges of any length- N accepting walk, and use the single central fragment. Such an accepting walk exists because the hull slice is nonempty. If $1 \leq N \leq 2e$, apply the lemma to the whole walk; its bound is no larger than (115). The empty word is harmless. Boundary and central fragments may have different witnessing seeds. $\square$

85.2 High exponent structure and the quadratic bound

Put $\omega_n = (n-1)^2 + 1$. We use the following classical facts from Kirkland, Olesky and van den Driessche [30]: Theorem 1.1, Lemma 2.1, Theorems 2.2–2.3, and Corollary 2.4, printed pages 30–34. Their Theorem 1.1 credits the two-cycle-length result to Lewin and Vitek.

For a primitive Hamiltonian graph with $n \geq 3$ and

$$e \geq \lfloor \omega_n/2 \rfloor + 2, \quad (116)$$

there is a unique Hamilton cycle; the other elementary cycles have one length j , where $\gcd(n, j) = 1$ and $n > j \geq \lceil (n-1)/2 \rceil$. If $j > n/2$, label vertices $1, \dots, n$ so the Hamilton edges descend modulo n , and the other edges are

$$i \longrightarrow i + j - 1 \quad (i \text{ in a subset of } \{1, \dots, a\}), \quad 1 \leq a \leq n - j + 1, \quad (117)$$

with sources $1, a$ present. Their exact exponent formula is

$$e = n - a + 1 + (n - 2)j. \quad (118)$$

The exceptional possibility is odd n with $j = (n-1)/2$; the same formula holds with $1 \leq a \leq j$. These imported statements concern the underlying graph, without labels.

Proof of Theorem 85.1. First assume $n \geq 4$, (116), and $j > n/2$. Set $c = j - a \geq 0$. Consider first $c \geq 1$ and the descending path

$$H: \quad j \longrightarrow j - 1 \longrightarrow \dots \longrightarrow a,$$

of c edges. Every elementary j -cycle uses one chord and contains H . To verify the one-chord assertion, a j -cycle using r chords has total displacement $-j + rj$ modulo n , so $r \equiv 1 \pmod{n}$. Since $0 \leq r \leq a < n$, it follows that $r = 1$. Its remaining edges descend from $i + j - 1$ back to its chord source i , and include all edges of H . The Hamilton cycle contains H as well.

No internal vertex of H has an outside entering or leaving edge. Indeed, chord sources are at most a , while chord targets lie between j and $a + j - 1 \leq n$. There is no modular wrap in the chord targets.

Every edge of H is doubly labeled. For either letter x , factor universality supplies arbitrarily long x -labeled walks. A sufficiently long such walk contains an elementary directed cycle all of whose edges carry x . Every elementary cycle contains H , so all edges of H carry x .

Use the weights of Lemma 85.3. Each j -cycle has total weight $a + 1$, and the Hamilton cycle has total weight $n - j + a + 1$. Their means obey

$$j(n - j + a + 1) - n(a + 1) = (n - j)(c - 1) \geq 0.$$

Thus the maximum cycle mean is

$$\mu = \frac{n - j + a + 1}{n} \leq 1.$$

By (115) and (118),

$$R(L) \leq \frac{2}{n}(n - j + a + 1)((n - 2)j - a + 2) + 2n + 1. \quad (119)$$

Put $k = n - j$ and $G(a) = (k + a + 1)((n - 2)j - a + 2)$. On the real interval $1 \leq a \leq k + 1$,

$$G'(a) = (n - 2)j - k + 1 - 2a \geq (n + 1)j - 3n - 1 \geq 0.$$

The last inequality follows from $j > n/2$ when $n \geq 6$, and from the integer bound $j \geq 3$ when $n = 4, 5$. Hence $G(a) \leq G(k+1)$. Substitution into (119) gives

$$\begin{aligned} R(L) &\leq \frac{4(n-1)}{n}(n-j+1)(j-1) + 2n + 1 \\ &\leq n(n-1) + 2n + 1 = n^2 + n + 1, \end{aligned}$$

since $(n-j+1) + (j-1) = n$. The endpoint $a = k+1$ is only an algebraic maximizer; it need not satisfy $c \geq 1$. The path lemma was applied at the actual graph's parameters.

If $c = 0$, then $a = j \leq n-j+1$ and $j > n/2$ force odd n and $a = j = (n+1)/2$. Equation (118) gives $e = (n^2 - 1)/2$, so Lemma 85.2 gives $R(L) \leq n^2$.

If (116) fails, then

$$R(L) \leq 2e + 1 \leq (n-1)^2 + 4 \leq n^2 + n + 1 \quad (n \geq 2).$$

If the exceptional case $j = (n-1)/2$ holds, (118) and $a \geq 1$ give

$$R(L) \leq 2e + 1 \leq n^2 - n + 3 \leq n^2 + n + 1.$$

These cover all $n \geq 4$ cases.

For $n = 1$, factor universality forces both self-loop labels, and $R(L) = 1$. For $n = 2, 3$, the classical primitive exponent bound $e \leq (n-1)^2 + 1$ gives $R(L) \leq 5, 11$, respectively. That bound and these small cases are recalled in [30, Section 1]. They satisfy the claimed estimate. $\square$

86 Graphs with a cycle on all but one vertex

The *circumference* of a digraph is the maximum length of an elementary directed cycle. The compression argument extends to the case where at most one vertex is omitted.

Theorem 86.1. *Let L be recognized by an n -state, ε -free, factor-universal binary NFA with primitive underlying digraph and circumference at least $n-1$. For arbitrary nonempty initial and final sets,*

$$R(L) \leq n^2 + 3n.$$

Lemma 86.2 (Several common paths). *In a primitive factor-universal n -state NFA with nonempty initial and final sets, suppose there are $t \geq 1$ pairwise vertex-disjoint directed paths of positive lengths, with C edges in total. Each edge carries both letters, internal path vertices have no outside incoming or outgoing edges, and every elementary cycle contains all the paths. If all cycle lengths are at most k , their mean weights under the first-edge charge are at most*

$$\mu = 1 - \frac{C-t}{k}.$$

The segment estimate (113) and global bound (115) hold with this μ , or with any larger valid cycle-mean bound in $[0, 1]$.

Proof. Charge every outside edge and the first edge of every path by one, and all other path edges by zero. A cycle of length ℓ has weight $\ell - (C-t)$, hence mean at most $1 - (C-t)/k$. Every path run except possibly an initial truncated run contains its charged first edge. Therefore pieces are bounded by weight plus one, just as in Lemma 85.3. The same residual simple-path estimate proves the claim. $\square$

Proof of Theorem 86.1. The Hamiltonian case follows from Theorem 85.1. Assume first $n \geq 6$, exponent satisfying (116), and circumference $k = n - 1$. By [30, Theorems 1.1–1.2], the elementary cycle lengths are k and j , with

$$\gcd(j, k) = 1, \quad j \geq \lceil n/2 \rceil, \quad e \leq n + j(n - 3). \quad (120)$$

If $j = n/2$, Lemma 85.2 gives $R(L) \leq n^2 - n + 1$, so assume $j > n/2$.

The classification [30, Theorem 3.3] places the graph, up to relabeling and reversal of all arcs, inside one of two templates. Start with the descending cycle on $1, \dots, k$ and every chord

$$i \rightarrow i + j - 1 \quad (1 \leq i \leq a_0 = n - j).$$

Either add vertex n with the same base predecessors and successors as one base vertex v , or add edges $1 \rightarrow n$, $n \rightarrow n - 2$, $n \rightarrow j - 1$.

Reversal causes no loss: reverse all labels' edges, exchange I and F , and reverse each word. This bijects seeds and hull words, preserves every aligned partition after reversing its interval order, and preserves factor universality. Thus $R(L)$ is unchanged.

In the base template, every elementary cycle contains the descending path H from j to a_0 , of length

$$c = j - a_0 = 2j - n.$$

The base cycles have lengths j and k , with a unique k -cycle and one chord in each j -cycle, by the Hamiltonian classification or the displacement argument used above. Internal vertices of H have unique incoming and outgoing base edges.

In the replication template, project n to v . Every elementary cycle projects to an elementary base cycle. Otherwise the only repeated projected vertex is v , and splitting there gives two nonempty closed base walks. Each has length at least the base girth j , forcing an original elementary cycle of length at least $2j > n$, which is impossible. Remove from H its at most two edges incident with v . The remaining edges form at most two vertex-disjoint nonempty paths after empty pieces are discarded. Every template cycle contains them, their internal vertices retain unique incoming and outgoing edges, and their total length C and number t satisfy

$$C \geq c - 2, \quad t \leq 2, \quad C - t \geq c - 4.$$

In the other template, the descending path from $j - 1$ to a_0 is common to all cycles and has $c - 1$ edges. For a cycle through n , entrance is from 1. Exit to $n - 2$ gives a descending return path and a k -cycle; exit to $j - 1$ gives a j -cycle. A chord on either return route would force a repeated base vertex before returning to 1, so no further possibility occurs. Internal vertices of this shortened path have unique incoming and outgoing edges. For $c \geq 2$ it is a positive-length path with $C - t = c - 2 \geq c - 4$. At $c = 1$ it is empty, a case handled below without compression.

Each retained edge lies in every cycle of the relevant full template, so it lies in every cycle of the actual strongly connected subgraph. In particular it is present there. Factor universality forces it to carry both letters, using a monochromatic elementary cycle for each letter as in the Hamiltonian proof. Internal degree conditions persist under deletion of edges.

For $c \geq 4$, Lemma 86.2 therefore permits the larger mean bound

$$\mu_0 = 1 - \frac{c - 4}{n - 1} = \frac{2n - 2j + 3}{n - 1} \in [0, 1].$$

If no positive-length path remains when $c = 4$, charging every edge by one gives the same bound $\mu_0 = 1$ directly.

If $c \leq 4$, use (120) and $j \leq (n + 4)/2$ to obtain

$$R(L) \leq 2e + 1 \leq 2n + (n + 4)(n - 3) + 1 = n^2 + 3n - 11.$$

If $c \geq 4$, (115) and (120) give

$$R(L) \leq \frac{2(2n - 2j + 3)(j(n - 3) + 1)}{n - 1} + 2n + 1.$$

The derivative of the numerator product with respect to j is

$$(n - 3)(2n + 3 - 4j) - 2 < 0 \quad (j \geq (n + 4)/2).$$

Evaluate at $j = (n + 4)/2$ to get $R(L) \leq n^2 + 3n - 9$. Both estimates imply the theorem.

Below the high-exponent threshold, use $R(L) \leq (n - 1)^2 + 4$. For $2 \leq n \leq 5$, the primitive exponent bound gives $R(L) \leq 2(n - 1)^2 + 3 \leq n^2 + 3n$. For $n = 1$ the rank is one, and an accepted empty word always has rank one. This exhausts the cases. $\square$

87 A large exponent without a total labeled state subset

A nonempty $S \subseteq Q$ is *forward-total* if every $q \in S$ has, for each letter, an outgoing edge into S . It is *backward-total* if every $q \in S$ has, for each letter, an incoming edge from S . The next family shows why such subsets cannot be assumed even near the maximum primitive exponent.

Proposition 87.1. *For every odd $n \geq 5$ there is a primitive, Hamiltonian, factor-universal binary graph on n vertices with exponent $(n - 1)(n - 2)$, no nonempty forward-total or backward-total subset, and*

$$R(L) \leq 6n + 3$$

for every choice of nonempty initial and final sets.

Proof. Use $Q = \{0, \dots, n - 1\}$, Hamilton edges $i \rightarrow i - 1$ modulo n , and chords

$$0 \rightarrow n - 3, \quad 2 \rightarrow n - 1.$$

The edges $0 \rightarrow n - 1$ and $1 \rightarrow 0$ carry only 0; the edge $n - 1 \rightarrow n - 2$ carries only 1. Every other edge, including both chords, carries both letters.

Set $A = Q \setminus \{n - 2\}$ and $C = Q \setminus \{0\}$. Directly,

$$\delta(S, 0) = A, \quad \delta(S, 1) = C \quad (S \in \{Q, A, C\}). \quad (121)$$

For completeness, deleting 0 can remove only targets $n - 1, n - 3$; alternative suppliers $2, n - 2$ remain in C . Deleting $n - 2$ can remove only target $n - 3$; its alternative supplier 0 remains in A . Together with the displayed full-set images this proves (121), letter by letter. Thus every nonempty word maps Q to A or C , according to its last letter, proving factor universality.

A forward-total subset must omit 1, which has no outgoing 1, and $n - 1$, which has no outgoing 0. Both targets of 2 are then absent; the unique outgoing chain excludes $3, \dots, n - 2$, and finally both targets of 0 are absent. A backward-total subset must omit 0, which has no incoming 1, and $n - 2$, which has no incoming 0. The predecessors of $n - 3$ are then absent; this exclusion propagates down to 1. The two predecessors 0, 2 of $n - 1$ are also absent. Both kinds of subset are therefore empty.

The elementary cycle lengths are n and $n - 2$. An elementary cycle cannot use both chords: after the chord from 0, it meets 2 before returning to 0, and using the second chord would revisit a vertex; the reversed order is analogous. The graph is primitive since n is odd. In the notation of (117), after increasing the vertex names by one, $j = n - 2$ and $a = 3$. The graph belongs to the class of [30, Theorem 2.2], so Corollary 2.4 there gives

$$e = n - 2 + (n - 2)^2 = (n - 1)(n - 2).$$

Its girth is $g = n - 2$.

Choose any $f \in F$. There is a suffix length $b \leq n - 2 = g$ such that $|R_b| \geq 2$. Follow predecessors of f along the Hamilton cycle until reaching $n - 3$ or $n - 1$, then take one further predecessor step; these are precisely the two vertices of indegree two. This takes at most $n - 2 - f$ steps if $0 \leq f \leq n - 3$, at most two if $f = n - 2$, and one if $f = n - 1$. The singleton final set $\{f\}$ has two predecessors at that first branch, and monotonicity gives the assertion for F .

If $N > e + b$, the nonempty word $w[e : N - b]$ maps $Q = P_e$ onto a set missing just one state, by (121). That image meets R_b , so this entire central interval is admissible. Singleton boundaries would already give $R(L) \leq e + b + 1 \leq (n - 1)^2$.

For the stronger linear estimate, fix any length- N accepting underlying walk. Choose the target's symbol at every doubly labeled edge and the prescribed symbol at each other edge. This gives a seed that can disagree only at singly labeled edges. A given source vertex appears at most $\lceil \ell/g \rceil$ times among the sources of a walk segment of ℓ edges: two occurrences are separated by a positive closed walk, of length at least g . Only three source vertices, 0, 1, $n - 1$, have singly labeled edges. Hence a segment has at most $3\lceil \ell/g \rceil$ potential disagreements. A segment with d disagreements is covered by at most $2d + 1$ admissible pieces: matching runs use the constructed seed, and each disagreement uses its singleton hull witness.

If $1 \leq N \leq e + b \leq ng$, this gives at most $6n + 1$ pieces. If $N > e + b$, apply it to the first $e = (n - 1)g$ edges and the last $b \leq g$ edges, using at most $6(n - 1) + 1$ and 7 pieces, respectively. Add the central piece to obtain $6n + 3$. The empty-word convention completes the proof. $\square$

The family has $e/n^2 \rightarrow 1$ but $R(L)/n^2 \rightarrow 0$, uniformly over the endpoint sets. Thus near-extremal exponent alone does not produce a quadratic rank lower bound. It also disproves the proposed route of finding a nonempty total original-state subset in every high-exponent factor-universal graph. Its large reachable subsets and sparse singly labeled edges instead make the rank small.

88 Scope and reproducible checks

The all-size theorem establishes the leading coefficient for every fixed finite alphabet, quantitatively as $F_q(s) \leq s^2 + \max\{8, 2q - 4\}s$. Together with the alphabet-independent lower family it also covers $q = o(s)$. It does not establish a leading-one estimate uniform over unrestricted alphabet growth or the optimal additive term. The binary finite function is known here exactly only through four states. The sharper Hamiltonian and near-Hamiltonian estimates do not have matching lower constructions in those classes. A near-maximal primitive exponent can coexist with linear rank and no nonempty forward-total or backward-total state subset.

The evidence has three distinct roles. First, the periodic-obstruction, component, stripe, entry-exit, bubble, circulation, common-anchor, parser, cutoff, pruning, and explicit-family arguments are ordinary proofs supplied in this article. The primitive exponent and high-exponent structural inputs are classical results cited precisely in [30]. The length-gate construction is reproduced from the companion note (Part IV). Second, the exact binary two-, three-, and four-state upper extrema depend on complete finite computations whose coverage is justified in Section 84. Independent implementations reproduce the counts, with pointwise

agreement for all two- and three-state presentations and a full independently regenerated pruned four-state run. Third, the finite parameter, bookkeeping, template, and family checks corroborate the all-size proofs; they are not premises in place of those proofs.

The accompanying source package contains the editable article, PDF, build script, exact implementations, reproduction scripts, compact count and audit receipts, consistency checks, pinned source descriptions, and checksums. Millions of raw automaton records and compiled executables are omitted. Candidate lists can be regenerated by the supplied programs, and their digests compared with the receipts. The README distinguishes fast receipt and consistency validation from complete exhaustive reruns.

The mathematical arguments and computation packages have undergone separate independent review, followed by review of the integrated article. This is an unrefereed research note with ordinary and computer-assisted proofs, not a Lean development or another proof-assistant certificate. The provenance claim is local: the results answer questions developed from the pinned report, without claiming first discovery throughout the literature. [write, batch 73] The reviews mentioned above are recorded in three delivered receipts (review-receipt.json, finite_alphabet_integration_audit.json and linear-exact-census-integration-review.json, shipped as listed below); they belong to the delivery and are not refereeing. [write, batch 73] No statement of this Part is formalized; Part IV's Question 53.4 applies to it as well. The exact $F_2(s)$ for $s \geq 5$, the optimal additive term and a leading constant uniform in the alphabet size remain open; they refine Question 53.1.

[write, batch 73] *Shipped layout.* The paragraphs above describe the delivered package. In this report its files are shipped, byte-identical, as follows; the manuscript and its PDF (delivered as hamiltonian-rank.tex and hamiltonian-rank.pdf, a file name kept from an earlier delivery) are printed or replaced by this report, the delivery README by the report's README, and the checksum list SHA256SUMS was verified and retired. Two delivered files are not shipped again because they are already shipped with Part IV, byte for byte: check_lower.py is code/04-quadratic-finite-rank-check.py and lower-verification.json is data/04-quadratic-finite-rank-verification.json. In the table the shipped name is code/ or data/, then the prefix 07-sharp-constant-, then the second column.

In	Shipped name after the prefix	Delivered as
code	alphabet-check_fixed_alphabet_lemmas.py	alphabet/check_fixed_alphabet_lemmas.py
code	alphabet-check_ternary_bookkeeping.py	alphabet/check_ternary_bookkeeping.py
code	build.sh	build.sh
code	check.py	check.py
code	check_general.py	check_general.py
code	check_templates.py	check_templates.py
code	exact-audit_census.cpp	exact/audit_census.cpp
code	exact-audit_literal.py	exact/audit_literal.py
code	exact-compare_pointwise.cpp	exact/compare_pointwise.cpp
code	exact-four_state-audit_families.py	exact/four_state/audit_families.py
code	exact-four_state-audit_four.cpp	exact/four_state/audit_four.cpp
code	exact-four_state-census_four_pruned.cpp	exact/four_state/census_four_pruned.cpp
code	exact-four_state-check_literal_four.py	exact/four_state/check_literal_four.py
code	exact-four_state-four_state_prune_budget.cpp	exact/four_state/four_state_prune_budget.cpp
code	exact-four_state-reproduce.sh	exact/four_state/reproduce.sh
code	exact-four_state-validate_receipts.py	exact/four_state/validate_receipts.py
code	exact-producer-census.cpp	exact/producer/census.cpp
code	exact-producer-census_graph.cpp	exact/producer/census_graph.cpp
code	exact-producer-exact_rank.py	exact/producer/exact_rank.py
code	exact-reproduce.sh	exact/reproduce.sh
code	exact-validate_receipts.py	exact/validate_receipts.py
code	verify.sh	verify.sh
data	alphabet-fixed_alphabet_independent_audit_receipt.json	alphabet/fixed_alphabet_independent_audit_receipt.json

In	Shipped name after the prefix	Delivered as
data	alphabet-fixed_alphabet_lemmas_receipt.json	alphabet/fixed_alphabet_lemmas_receipt.json
data	alphabet-ternary_bookkeeping_receipt.json	alphabet/ternary_bookkeeping_receipt.json
data	exact-audit_receipt.json	exact/audit_receipt.json
data	exact-four_state-audit_receipt.json	exact/four_state/audit_receipt.json
data	exact-four_state-family_verification.json	exact/four_state/family_verification.json
data	exact-four_state-four-state-pruned-census.log	exact/four_state/four-state-pruned-census.log
data	exact-four_state-independent-pruned-census.log	exact/four_state/independent-pruned-census.log
data	exact-four_state-literal_rank_seven.json	exact/four_state/literal_rank_seven.json
data	exact-independent_one_state.json	exact/independent_one_state.json
data	exact-independent_three_state.json	exact/independent_three_state.json
data	exact-independent_two_state.json	exact/independent_two_state.json
data	exact-literal_receipt.json	exact/literal_receipt.json
data	exact-original_cutoff_rerun.json	exact/original_cutoff_rerun.json
data	exact-original_graph_rerun.json	exact/original_graph_rerun.json
data	exact-pointwise_three_state.json	exact/pointwise_three_state.json
data	exact-pointwise_two_state.json	exact/pointwise_two_state.json
data	finite_alphabet_integration_audit.json	finite_alphabet_integration_audit.json
data	general-verification.json	general-verification.json
data	linear-exact-census-integration-review.json	linear-exact-census-integration-review.json
data	review_receipt.json	review_receipt.json
data	source-pins.json	source-pins.json
data	verification.json	verification.json

The delivered `verify.sh` copies the checkers, receipts and the alphabet and exact directories into a temporary directory and runs them there, so it needs the delivered layout; the report's README gives a rerun on a copy that restores it. When the batch was placed, the ten Python checks that `verify.sh` runs without its `--compile` option were rerun on such a copy with Python 3.14.4; all passed, and those with a shipped record printed it apart from CRLF line endings. The two alphabet checkers also rewrite their receipts beside themselves, and those differ from the shipped ones only in the recorded running time (the "seconds" field). The C++ censuses and the two reproduction scripts were not rerun; their receipts validate.

A The classical universality input

[Added 30 September 2026, batch 66: Appendices A–D belong to Part I and are printed after Parts II and III so that no letter or number of Part I moved.] [Added 1 October 2026, batch 71: Part IV is printed before them too.] [Added 1 October 2026, batch 73: So are Parts V–VII.]

For completeness, here is the standard route to the co-r.e.-completeness used in Section 4. The foundational undecidable problem is Post correspondence [7]. Take two morphisms $h, g : I^* \rightarrow A^*$ and ask whether some $x \in I^+$ satisfies $h(x) = g(x)$. The set of positive instances is recursively enumerable, and the usual effective computation-history reduction gives r.e.-completeness.

Use disjoint alphabets I, A and a separator $\$$. Among encodings $x\$y^R$, where $x \in I^+$ and $y \in A^*$, the language

$$\{x\$y^R : y \neq h(x)\}$$

is context-free. A pushdown automaton pushes the successive output blocks of h while reading x , then compares the reversed candidate output against the stack. It accepts a symbol mismatch, a short candidate, or a long candidate; the regular encoding conditions can be imposed separately. The analogous language for g is also context-free. Their union, together with all malformed encodings, is universal exactly when the correspondence instance has no solution. This proves co-r.e.-hardness of grammar universality. Its complement is r.e. because grammar membership is decidable and nonmembers can be searched for by length.

To obtain a binary alphabet, map each letter of the finite encoding alphabet to a distinct binary block of one common length ℓ . For a language H over that alphabet, set

$$H' = c(H) \cup (\{0, 1\}^* \setminus c(A_0^*)),$$

where A_0 denotes the entire encoding alphabet. The image $c(H)$ is context-free and $c(A_0^*)$ is regular. The injective fixed-length code gives $H' = \{0, 1\}^*$ exactly when $H = A_0^*$. This is a reduction for universality; as Section 11 emphasizes, it is *not* automatically a reduction preserving crossover.

B Why the positional counts are eventually affine

The special case of Presburger counting needed here also has a direct one-dimensional explanation. After quantifier elimination, a predicate in variables (n, i) is a Boolean combination of linear inequalities and fixed modular congruences. We also impose $0 \leq i < n$.

Refine the residue class of n modulo all congruence moduli, and refine the residue class of i similarly. On these classes, congruences have constant truth values. Each nontrivial inequality in i changes truth value at a rational-affine threshold in n . The order of finitely many such thresholds is fixed once n is sufficiently large: two affine functions either agree identically or their order changes at most once.

Thus the allowed positions eventually form a fixed finite union of intervals between affine thresholds, possibly with selected endpoint positions, in each fixed residue class of i . Counting that class inside such an interval is a difference of floor or ceiling functions of rational-affine expressions. Refine the period of n to absorb the finitely many denominators. Each such count then is affine on a progression. Finite union and disjoint refinement preserve that property.

This argument explains both effectivity and the degree-one conclusion of Lemma 8.2; the general counting theorem of Woods supplies a systematic framework. Since counts are nonnegative for all sufficiently large parameters and integer-valued, their progression increments are nonnegative integers. That final observation is what makes the geometric factors in Theorem 8.3 polynomials rather than Laurent expressions.

C A compact proof-dependency map

Result	Essential inputs
Exact contraction and generation	Definition of aligned crossover; restriction and pairing of consecutive intervals. No grammar theorem.
One-step CFG undecidability	One-cut guard repair; classical fixed-binary CFG universality.
Eventual CFG undecidability	Exact generation; disjoint boundary intervals for forbidden blocks; effective CFL union, concatenation, and star.
Regular-input stabilization	Exact generation; explicitly constructed distance automaton; classical limitedness decidability.
Rational weighted hull series	Effective Parikh semilinearity; one-parameter Presburger counting; coordinate-product enumeration.
Non-CFL linear-seed example	Explicit grammar; elementary position inequalities; CFL pumping lemma and intersection with a regular language.
Minimal finite-stage recurrence	Exact binomial count; degree of a polynomial; reduced rational denominators.

D [write] Provenance of this report

One manuscript. This report is built from a single manuscript, manuscript 02 of batch 44 of ProveIt’s incoming-report intake: *Aligned Fragments and Constrained Crossover*, dated 29 September 2026, delivered as the archive `ProveIt_Constrained_Crossover.zip`, which arrived in commit `ae28ea2db` and was placed in commit `203016015`. Its pin is ProveIt commit `9b24a3a8d545af9624f6ac455f5b548be62818b6` (Section 1). The manuscript contributed everything outside the text marked [write]: every result, proof, example, remark, question, limitation and the Lean plan, all printed. The package’s audit notes and verifier are shipped unchanged beside this article (Section 10.2 lists the paths). Its 23-page US-letter PDF and its README are not shipped; they survive in the arrival commit. [Added 30 September 2026, batch 66: this describes Part I; Parts II and III come from two further manuscripts (last paragraph of this appendix).]

What the write added. Everything marked [write]: the title-page note; in Section 1, where Hughes poses the question, the state of the sibling source, and Sections 1.4 and 1.5; the note after Remark 2.3; the notes on the shipped layout and on formal status in Section 10; the note on Hughes’s Sections 13–14 after the research questions; and this appendix. The locations in Hughes’s preprint were checked against its arXiv v1 PDF when the report was written.

Choices. There was no merge, so no result had to be chosen over another. Every label received the prefix `ccc:`, and every reference was updated with it. Labels were added only on new text and on Question 11.1 and Section 11.1; no theorem, section or equation number of the manuscript changed. The page size stays US letter, as delivered. No symbol was renamed; the `shuffle` package was added for the symbol $\sqcup$ of the sibling report.

[Added 30 September 2026, batch 66: Two further manuscripts.] This appendix describes the report as written in batch 44, which is Part I. On 30 September 2026 two manuscripts of batch 66, both written against Part I at ProveIt commit `b8b0fa218` (where Part I’s text was the one printed here), were added as Parts II and III: manuscript 01, *Finite Crossover Stabilization Is PSPACE-Complete* (archive

ProveIt_Crossover_Stabilization.zip), and manuscript 03, *Decidable Regularity and Exact Stabilization of Crossover Closures* (archive ProveIt_Crossover_Classification.zip). Both arrived in commit 9ad899cbe and were placed in commit 4fee1cd07. The report is therefore built from three manuscripts. Sections 13.1 and 29.1 give each one’s provenance and what it contributed. The choices made were these. The two were kept as separate Parts, not merged, because they share no main theorem; the only overlap is a reproof of Part I’s coordinate hull, printed once (Theorem 3.3) with the two reproofs kept as marked second and third routes (Proposition 15.2, Lemma 31.3). Part III’s Lemma 31.4, which restates Part I’s Lemma 8.1 without citing it, now credits it, and its proof is kept as a second route. Every label of the two manuscripts received the prefix `ccc:ps:` or `ccc:cl:`; their appendices became sections, because the appendices belong to Part I; ten symbols were renamed where they clash (Sections 13.3 and 29.3); the bibliographies were merged (each manuscript’s entry for Part I replaced by internal references, its Hughes entry identified with Part I’s, and Part III’s Esparza et al. entry with [3]). Four labels were added to existing text of Part I (`ccc:q:regular`, `ccc:q:grammar`, `ccc:q:series`, `ccc:q:restricted` on its questions) and dated notes were added after Corollary 7.4, after four of its questions, in its Sections 10–11, in its conclusion, on the title page, and here; no existing label, theorem, section, equation or question number of Part I changed.

[Added 1 October 2026, batch 71: A fourth manuscript.] On 1 October 2026 batch-71 manuscript 01, *Quadratic Finite Rank for Constrained Crossover* (archive quadratic-finite-rank-package.zip), written against Parts I–III at ProveIt commit 1bd730779, which arrived in commit aee32ad45 and was placed in commit 427bca743, became Part IV. The report is now built from four manuscripts; Section 46.1 gives this one’s provenance and what it contributed. The choices made were these. Part IV answers Part II’s Question 24.2 and imports Part II’s upper bound; it was kept as a Part of its own, not merged into Part II, because its construction is new and shares no proof with Part II. Its restatement of the rank calculus is printed as a marked restatement with its proof as one more route (Proposition 48.1); its last-gap lemma received a credit to Sylvester’s two-generator Frobenius formula; its gate language was renamed L_m^{gate} , because Part II’s L_m is another family; its four unnumbered questions received numbers and labels; every label received the prefix `ccc:qr:`; its bibliography entry for the pinned report was replaced by internal references, its To entry identified with [8], and two entries (Sylvester, Ramírez Alfonsín) were added for the credit. Dated notes were added in Part I’s Section 10, its conclusion and Appendix A, in Part II’s delivered status, after Theorem 22.1, after Questions 24.2 and 24.8, in Part II’s conclusion, and here (the title page, already full, received none); no existing label, theorem, section, equation or question number changed.

[Added 1 October 2026, batch 73: A fifth manuscript.] On 1 October 2026 batch-73 manuscript 33, *Rational Rank Slopes for Constrained Crossover* (archive rational-rank-slopes-package.zip), written against Parts I–III at ProveIt commit 1bd730779, which arrived in commit f8c3a392a and was placed in commit 5e4f1263c, became Part V. The report is now built from five manuscripts; Section 55.1 gives this one’s provenance and what it contributed. The choices made were these. Part V answers Part II’s Question 24.3 and was kept as a Part of its own, because its theorem and its proof (greedy optimality, the subset-reset graph, cycle means) are new; its restatement of Part II’s aligned fragment criterion became a pointer without the repeated proof, its periodic definitions and its parity example are kept with pointers to Part II; its slope λ_r was renamed σ_r , because Part III’s λ_v is another rational quantity; its sentence of next questions became four numbered questions; every label received the prefix `ccc:rs:`; its bibliography entries for the pinned report and for Part IV’s manuscript were replaced by internal references, its Manzoni entry identified with [19], and its Karp (1978) entry printed as [26], because [17] is another paper. Dated notes were added after Questions 24.2, 24.3, 53.1 and 53.2, after Part II’s fixed-word algorithm and Corollary 18.5, in Part II’s conclusion, in Part I’s conclusion, in Part IV’s Section 53, in Appendix A and here; no existing

label, theorem, section, equation or question number changed.

[Added 1 October 2026, batch 73: A sixth manuscript.] On 1 October 2026 batch-73 manuscript 29, *Finite Crossover Stabilization for Deterministic Automata* (archive `dfa-crossover-package.zip`, revised edition), written against Parts I–III at ProveIt commit 1bd730779, which arrived in commit f8c3a392a and was placed in commit 26abf259b, became Part VI. Its earlier edition, manuscript 32 (`dfa-crossover-package (1).zip`), was superseded and is neither printed nor shipped. The report is now built from six manuscripts; Section 66.1 gives this one’s provenance and what it contributed. The choices made were these. Part VI answers Part II’s Question 24.1 and Part IV’s Question 53.2 and was kept as a Part of its own. Its two lemmas on actual alphabets and on the joint entry and exit budget are printed here, with their proofs, once; a later manuscript states them for NFAs with the same proofs. Its restatement of the report’s background is kept with pointers; its gate lemma, whose first assertions are Part IV’s, keeps its proof, because its labelling and its common completion differ. Its language L_m was renamed L_m^{dfa} and the initial state r of its gate automaton q_{in} ; a lost $\text{\rm}$ in one display was restored; its sentence of open refinements became four numbered questions; every label received the prefix `ccc:dc:`; its bibliography entries for the pinned report and for the manuscripts of Parts IV and V were replaced by internal references, its Schwarz entry identified with [27], and two entries (Stockmeyer–Meyer, Gawrychowski et al.) added. Dated notes were added wherever the report called the DFA complexity or the DFA order open: in Part I after Corollary 7.4 and Question 11.2 and in its conclusion; in Part II’s opening, delivered status, comparison table, Section 21, the comment after Theorem 22.1, after Questions 24.1, 24.2 and 24.5, and in its conclusion; in Part III after its complexity remark, after Question 42.3 and in its conclusion; in Part IV’s opening, its list of what it leaves unaffected, after Questions 53.1 and 53.2 and after its remark on the DFA decision question; in Appendix A and here. No existing label, theorem, section, equation or question number changed.

[Added 1 October 2026, batch 73: A seventh manuscript.] On 1 October 2026 batch-73 manuscript 21, *Sharp Leading Constant for Constrained Crossover Rank* (archive `common-path-rank-package.zip`), written against Parts I–III at ProveIt commit 1bd730779 and citing Part IV’s manuscript at commit aee32ad45, which arrived in commit f8c3a392a and was placed in commit 5e4f1263c, became Part VII. Its earlier version, manuscript 24 (`common-path-rank-package (1).zip`), was superseded and is neither printed nor shipped. The report is now built from seven manuscripts; Section 75.1 gives this one’s provenance and what it contributed. The choices made were these. Part VII answers Part IV’s Question 53.1 as to the leading constant and was kept as a Part of its own. It shares two lemmas, with the same proofs, with Part VI; they are printed once, with their proofs, in Part VI, and Part VII states them for NFAs and points there. Its reproofs of Part II’s criterion and transient bound and of Part IV’s construction became pointers without the repeated proofs; its greedy lemma, Part V’s, keeps its proof as a second proof. Its backward layers B_j were renamed R_j , Part II’s name; four sections without labels received labels; every label received the prefix `ccc:sc:`; its bibliography entries for the pinned report and for the manuscripts of Parts IV and V were replaced by internal references, its Schwarz entry (not cited in its text) identified with [27], and one entry (Kirkland–Olesky–van den Driessche) added. Two of its delivered files are Part IV’s shipped checker and record and were not shipped again. Dated notes were added wherever the report called the leading constant open or bounded it by two: after Questions 24.2 and 53.1, after the comment on Theorem 22.1, in Part II’s delivered status and conclusion, in Part IV’s opening, its list of answers, its single-initial-state paragraph and Section 53, in Part V after Corollary 64.2 and in its list of what it leaves open, and in Part VI after Theorem 73.1, after its two shared lemmas, at the end of Section 73 and in its list of what it leaves open; also in Part I’s conclusion, in Appendix A and here. No existing label, theorem, section, equation or question number changed.

References

- [1] Charles E. Hughes. *Undecidability of Adjacent Equality for Insertion, Shuffle, and Crossover Language Operations*. arXiv:2608.27755v1, 27 August 2026. Definition on printed p. 2; the constrained one-step question on printed p. 3. [write] Section 2 (definition), Section 4 (the question, pp. 3–4), Section 14 (open directions, p. 12). <https://arxiv.org/abs/2608.27755v1>.
- [2] Vladimir Reshetnikov and ProveIt contributors. *ProveIt: Gaps in Bounded-Shuffle Hierarchies*, Parts I–II. Inspected 29 September 2026 at the revision and path specified in Section 1. AI-assisted research drafts. [write] The sibling report *insertion-degree-spectra* of this collection (Section 1.4). <https://github.com/VladimirReshetnikov/ProveIt>.
- [3] Javier Esparza, Pierre Ganty, Stefan Kiefer, and Michael Luttenberger. Parikh’s theorem: A simple and direct automaton construction. *Information Processing Letters* **111**(12) (2011), 614–619. <https://arxiv.org/abs/1006.3825>. doi:10.1016/j.ipl.2011.03.019.
- [4] Kevin Woods. Presburger arithmetic, rational generating functions, and quasi-polynomials. *Journal of Symbolic Logic* **80** (2015), 433–449. <https://arxiv.org/abs/1211.0020>. doi:10.1017/jsl.2015.4.
- [5] Daniel Kirsten. Distance desert automata and the star height problem. *RAIRO—Theoretical Informatics and Applications* **39**(3) (2005), 455–509. https://www.numdam.org/item/ITA_2005__39_3_455_0/. doi:10.1051/ita:2005027.
- [6] Markus Krötzsch, Tomáš Masopust, and Michaël Thomazo. Complexity of universality and related problems for partially ordered NFAs. *Information and Computation* **255**, Part 1 (2017), 177–192. Preprint arXiv:1609.03460v2. <https://arxiv.org/abs/1609.03460>.
- [7] Emil L. Post. A variant of a recursively unsolvable problem. *Bulletin of the American Mathematical Society* **52** (1946), 264–268.
- [8] Anthony Widjaja To. *Unary finite automata vs. arithmetic progressions*. *Information Processing Letters* 109(17) (2009), 1010–1014. Corrected proof in Section 3 of arXiv:0812.1291v1. <https://arxiv.org/abs/0812.1291>. DOI: 10.1016/j.ipl.2009.06.005.
- [9] Jui-Yi Kao, Narad Rampersad, and Jeffrey Shallit. *On NFAs Where All States are Final, Initial, or Both*. *Theoretical Computer Science* 410(47–49) (2009), 5010–5021. Lemma 6 in arXiv:0808.2417v2, revised 3 July 2009. <https://arxiv.org/abs/0808.2417>. DOI: 10.1016/j.tcs.2009.07.049.
- [10] Walter J. Savitch. *Relationships between nondeterministic and deterministic tape complexities*. *Journal of Computer and System Sciences* 4(2) (1970), 177–192. DOI: 10.1016/S0022-0000(70)80006-X.
- [11] Rohit J. Parikh. On context-free languages. *Journal of the ACM* **13**(4) (1966), 570–581. doi:10.1145/321356.321364.
- [12] Seymour Ginsburg and Edwin H. Spanier. Semigroups, Presburger formulas, and languages. *Pacific Journal of Mathematics* **16**(2) (1966), 285–296. doi:10.2140/pjm.1966.16.285.
- [13] Pablo Barceló, Chih-Duo Hong, Xuan-Bach Le, Anthony W. Lin, and Reino Niskanen. Monadic decomposability of regular relations. In *ICALP 2019, LIPIcs* **132**, article 103. doi:10.4230/LIPIcs.ICALP.2019.103.
- [14] Matthew Hague, Anthony W. Lin, Philipp Rümmer, and Zhilin Wu. Monadic decomposition in integer linear arithmetic. In *IJCAR 2020, Part I, LNCS* **12166**, 122–140. doi:10.1007/978-3-030-51074-9_8. Extended version: <https://arxiv.org/abs/2004.12371>.
- [15] Moses Ganardi and Marin Ricos. Constructing small monadic decompositions in Presburger arithmetic. In *LICS 2026, LIPIcs* **380**, 47:1–47:22. Published 9 July 2026. doi:10.4230/LIPIcs.LICS.2026.47.
- [16] Danny Raz. Length considerations in context-free languages. *Theoretical Computer Science* **183**(1) (1997), 21–32. doi:10.1016/S0304-3975(96)00308-8.
- [17] Richard M. Karp. Reducibility among combinatorial problems. In R. E. Miller and J. W. Thatcher (eds.), *Complexity of Computer Computations*, Plenum, 1972, 85–103. doi:10.1007/978-1-4684-2001-2_9.
- [18] J. Barkley Rosser and Lowell Schoenfeld. Approximate formulas for some functions of prime numbers. *Illinois Journal of Mathematics* **6**(1) (1962), 64–94. doi:10.1215/ijm/1255631807.
- [19] Luca Manzoni, Leonardo Vanneschi, and Giancarlo Mauri. *A distance between populations for one-point crossover in genetic algorithms*. *Theoretical Computer Science* 429 (2012), 213–221. DOI: 10.1016/j.tcs.2011.12.041. Author-repository full text: <https://run.unl.pt/server/api/core/bitstreams/1639e177-ebcc-4aa6-baf2-ab9778007dd2/content>.

- [20] Helmut Wielandt. *Unzerlegbare, nicht negative Matrizen*. *Mathematische Zeitschrift* 52 (1950), 642–648. DOI: [10.1007/BF02230720](https://doi.org/10.1007/BF02230720).
- [21] Stewart W. Neufeld. *A diameter bound on the exponent of a primitive directed graph*. *Linear Algebra and its Applications* 245 (1996), 27–47. DOI: [10.1016/0024-3795\(94\)00203-7](https://doi.org/10.1016/0024-3795(94)00203-7).
- [22] J. J. Sylvester. *Mathematical questions with their solutions*. *Educational Times* 41 (1884), 21.
- [23] Jorge L. Ramírez Alfonsín. *The Diophantine Frobenius Problem*. Oxford Lecture Series in Mathematics and its Applications 30, Oxford University Press, 2005.
- [24] Maxime Crochemore, Alessio Langiu, and Filippo Mignosi. Note on the greedy parsing optimality for dictionary-based text compression. *Theoretical Computer Science* 525 (2014), 55–59. doi:[10.1016/j.tcs.2014.01.013](https://doi.org/10.1016/j.tcs.2014.01.013). arXiv:[1211.5350](https://arxiv.org/abs/1211.5350).
- [25] Alessandro De Luca and Gabriele Fici. Dorst–Smeulders coding for arbitrary binary words. arXiv:[2509.08684v1](https://arxiv.org/abs/2509.08684) (10 September 2025).
- [26] Richard M. Karp. A characterization of the minimum cycle mean in a digraph. *Discrete Mathematics* 23(3) (1978), 309–311. doi:[10.1016/0012-365X\(78\)90011-0](https://doi.org/10.1016/0012-365X(78)90011-0).
- [27] Štefan Schwarz. On the semigroup of binary relations on a finite set. *Czechoslovak Mathematical Journal* 20(95) (1970), no. 4, 632–679. Theorem 4.3, p. 653; index definition, p. 634. doi:[10.21136/CMJ.1970.100989](https://doi.org/10.21136/CMJ.1970.100989). Primary full text, theorem page.
- [28] Larry J. Stockmeyer and Albert R. Meyer. Word problems requiring exponential time (preliminary report). *Proceedings of the Fifth Annual ACM Symposium on Theory of Computing* (1973), 1–9. Theorem 6.1, pp. 7–8. doi:[10.1145/800125.804029](https://doi.org/10.1145/800125.804029). Author-hosted original text.
- [29] Paweł Gawrychowski, Martin Lange, Narad Rampersad, Jeffrey Shallit, and Marek Szykuła. Existential length universality. arXiv:[1702.03961v5](https://arxiv.org/abs/1702.03961) (9 March 2020), full version, Theorems 5–6.
- [30] Steve Kirkland, D. D. Olesky and P. van den Driessche. Digraphs with large exponent. *Electronic Journal of Linear Algebra* 7 (2000), 30–40. Theorems 1.1–1.2, Lemma 2.1, Theorems 2.2–2.3 and 3.3, Corollary 2.4. Primary full text.