

Five-Particle Binary Automata

Reversible universality, a literal universal source,
startup and cellular clocks, exact lazy evaluation,
quartic certificates, a parallel replacement rule
and smaller sufficient radii

Research Reports 14–20, 26–28, 70 and 71 of one AI-assisted research pipeline
merged from twelve manuscripts (batch 82, clusters M1 and M2; batch 91, cluster W4)

3–4 October 2026

Abstract

One fixed kind of object runs through twelve research reports: a one-dimensional cellular automaton on the full binary shift that conserves the number of ones, is often globally reversible, and computes with exactly five ones. Part [I](#) proves that five particles suffice. Report 15 compiles every finite partial-injective guarded two-counter machine with a terminal halt into a globally reversible, number-conserving binary automaton of explicit finite radius, by an all-input isolated-swap lemma and a reflection argument; Report 14 compiles every deterministic two-counter program into a binary conservative automaton and instantiates it with a literal 8,408-row universal source (radius 756,787). With the separately proved four-mass decision theorem (Report 12), five is the exact particle threshold in both classes. Part [II](#) (Report 16) prints a literal universal reversible two-counter source, 122,622 controls and 141,561 rows, with a clean loader and all-input proofs; on its clean family positive return and periodicity of the fixed automaton are r.e.-complete with least period exactly $2\Theta + 2$. Part [III](#) (Reports 17 and 18) changes four history tags to remove an empty-startup cost of about $8 \cdot 10^{13}$ steps (138 executed steps remain), proves paired-boundary domination of the clocks and first-encounter optimality among 2^{233} static orientations, and lifts both to the automaton's clock through the identity $\Theta = \lambda M + \text{TV}(A^2) + \text{TV}(B^2) + U$. Part [IV](#) (Report 19) evaluates the frozen ordered product of 269,291,358,255 involutions exactly on every finite support without building it. Part [V](#) (Report 20) certifies that ordered rule by a degree-four sum of squares whose natural fiber is a singleton exactly when T steps complete within a budget of K changing factors, using $T + K$ rounds and no array over the factors. Part [VI](#) (Reports 26–28) replaces the ordered product by two parallel full-shift involutions that agree with it on every admissible state (sufficient radius $180D + 258 + 9J$, that is 91,711,698 for the universal source), evaluates the new rule exactly on every finite support with at most $n + 2$ discovery calls per step, and certifies it at each fixed horizon by a quartic with a unique natural witness. Part [VII](#) (Reports 70 and 71) gives two further full-shift extensions of the same admissible dynamics with smaller sufficient radii, by separating a short recognition radius from the contextual control radius and then shortening the triple exactness windows: $108D + 149 + 3J$ and $76D + 105 + 3J$ (55,027,013 and 38,722,713 for the universal source); each new rule differs from Report 26's on malformed inputs, and Report 70 adds that the only injective number-conserving binary rules of radius two are the five shifts. Both compilers of Part [I](#) come with fixed-horizon degree-four sum-of-squares certificates whose natural fibers are empty or singletons. These are AI-assisted research manuscripts with executable audits; they are unrefereed, not formalized, and make no novelty or priority claim. Re-proofs of results printed in neighbouring reports of the collection are kept as marked second presentations.

Contents

1	The series and this report	11
1.1	Twelve Reports on one automaton	11
1.2	Where the rest of the series is printed	12
1.3	How the Reports depend on one another	12
1.4	What the Parts prove	13
1.5	Questions answered inside the series and elsewhere	14
1.6	Relation to the collection and to the formal project	15
2	Conventions and notation	16
3	Status and non-claims	20
I	Five particles suffice in binary automata (Reports 15 and 14)	24
	<i>Report 15: Globally reversible binary computation with five particles</i>	24
4	Scope and theorem	24
5	Source machine and primary attribution	25
5.1	What is imported from Morita, and what is added here	26
6	An isolated swap on every binary input	26
7	Five-particle geometry	28
8	The partial micrograph and its exact clock	28
9	The edge matching as concrete pattern swaps	29
9.1	Free flight	29
9.2	Contact travel	29
9.3	Dispatch, endpoint and commit	30
9.4	Reading bounded counter classes without storing anything	30
10	Phase flip, reflection and observation	31
11	The exact resource ledger	32
12	Counted finite-horizon certificates	33
12.1	Disjoint finite class normalization	33
12.2	Every residual explicitly	34
12.3	Paying for the exact CA clock	34
12.4	Exact counts, expansion cost and height	35
12.5	The paid nonnegative-real variant	36
12.6	Materialized certificate of the small source	36
13	Clean exact targets without another particle	36
13.1	Full-history and projected clean certificates	38
14	A fully materialized accepting example	39
15	Verification, reproducibility and limitations	39
16	The five-particle threshold and relation to earlier work	42
17	Research questions and sharpenings	42
	<i>Report 14: Five particles in a binary conservative cellular automaton</i>	44
18	The model and the result	44

19	Finite codes and buffered geometry	45
20	The complete rule on arbitrary configurations	46
20.1	Free travel and departure	46
20.2	Home dispatch and the contextual zero test	46
20.3	Endpoint reversal and home commit	47
21	Conservation with remote reads and bounded writes	47
22	A finite local rule and its radius	48
23	The all-input simulation proof	49
23.1	Exact microstates of an arithmetic excursion	49
24	Exact clocks and the anchored halt word	50
25	Exact finite description counts	51
26	The pinned universal source	51
26.1	The finite-input primary machine	52
26.2	Finite half tapes and three natural counters	53
26.3	Three counters in two by prime coding	53
26.4	What the source verifier proves	54
26.5	The resulting fixed universal CA	55
27	The sharp particle threshold and its scope	55
28	A source horizon polynomial compiler	56
28.1	All residuals and the witness variables	57
28.2	Exact dimensions and algebraic costs	58
28.3	Adding exact physical time	58
28.4	Numerical height and space bounds	59
29	The paid nonnegative real variant	59
30	Literal and small fully emitted examples	60
30.1	A complete accepting witness and CA trace	61
30.2	External input variants and their costs	61
31	Executable specification and independent checks	63
31.1	Portable replay and artifact identity	64
32	Further research directions	64
II	A literal universal reversible source (Report 16)	66
	<i>Report 16: A literal universal reversible two counter source</i>	66
33	The machine and the exact claim	66
33.1	What is literal and what remains a finite specification	67
34	The finite tape and the irreversible front end	67
34.1	Primary table and exact finite loader	68
34.2	Three counters with the scanned symbol in control	69
34.3	Separate tests and a genuinely fresh start	70
34.4	Reducing indegree without changing data	70
35	Reversible history with five counters	71
35.1	Why history recording is globally reversible	72

36	Prime encoding into exactly two counters	72
36.1	Common transfer for arithmetic moves	73
36.2	Multiplication and an encoded increment	73
36.3	Division and an encoded decrement	73
36.4	Division with a bounded remainder for tests	74
36.5	Destination sharing is essential for restoration	74
36.6	Partial injection on every pair and strict serialization	75
37	Composition on all promised inputs and exact clocks	76
37.1	Finite prefixes and the no blocking theorem	76
37.2	The complete clock composition	76
37.3	Predicted clocks and what was actually executed	77
38	The fixed binary five particle instance	77
38.1	Indexed construction and global reversibility	78
38.2	All compiler parameters for this source	78
38.3	Loader and exact CA observation clock	79
38.4	Positive return and periodicity now have an unconditional clean family theorem	80
38.5	Relationship to the four mass lower bound and exact targets	81
39	Exact finite horizon quartic certificate counts	81
39.1	Disjoint class normalization and counted witnesses	81
39.2	Every residual and the full natural zero fiber	82
39.3	Exact cost and optional clocks or real domain	83
39.4	A smaller certificate using shared primitive test offsets	84
39.5	The improved counts and the price of expansion	85
39.6	Certificates for least period and return multiples	86
40	Evidence and portable reproduction	87
40.1	Independent layers of evidence	87
40.2	Checker hardening before the final freeze	88
40.3	Release inventory and exact manifest	88
40.4	One fresh standard library replay	88
40.5	What remains unmaterialized	89
41	Further questions that follow from this release	89
III	Startup histories and cellular clocks (Reports 17 and 18)	91
	<i>Report 17: Startup history in a literal reversible source</i>	91
42	The machine and the precise comparison	91
43	The exact four swaps	92
44	Why the whole original input theorem survives	93
44.1	The history invariant is independent of the bit assignment	93
44.2	Prime graphs and exact literal costs	94
44.3	Global injection and nonblocking simulation are separate	94
45	Exact initialization for all natural scratch values	95
45.1	Zero scratch keeps the entire history at zero	95
45.2	No nonzero-scratch input is removed	95
46	The order preserving comparison proof	96
47	The exact optimum within the orientation class	98
48	The genuine equal history slowdown	99

49	The executed 348 step prefix	99
50	The binary compiler and the limits of the clock claim	101
51	Reproducibility and provenance	101
52	What is resolved and what remains useful	103
	<i>Report 18: Cellular clocks for a literal reversible source</i>	104
53	The comparison and its domain	104
54	An exact total variation identity	106
55	Whole prime macros from their literal paths	107
55.1	Increment and enabled decrement	107
55.2	Successful zero and positive tests	107
56	Monotonicity and the common compiler constant	108
57	The recorder embedding charges complete CA macros	109
58	Prefix optimality and the four startup swaps	110
59	The exact minimum when initial scratch is zero	111
60	An exact formula for the recorded suffix	112
61	What was executed and what was inferred	113
62	Connection to finite horizon timing certificates	114
63	Reproducibility and scientific provenance	115
64	Conclusions and remaining boundaries	116
IV	Exact lazy evaluation (Report 19)	117
	<i>Report 19: Exact lazy finite support evaluation of the frozen reversible binary CA</i>	117
65	Result and scope	117
66	Pinned source semantics and metadata	118
66.1	Input schema and validation	118
66.2	Constants and stored data	118
67	Arithmetic access to the ordered factors	119
68	Complete discovery of possible changes	120
69	Exact sparse application of a factor	121
69.1	Raw keys and eligibility	121
69.2	Simultaneous substitution	121
70	Ordered skipping and exact inverse	122
71	A malformed input with a real cascade	122
72	Complexity and resource limits	123
72.1	Source metadata and class cutoff	123
72.2	Expected hash model	124

72.3	Adversarial hashing and exact integers	124
73	Executable evidence and universal observations	124
73.1	Differential regression and independent audit	124
73.2	Literal universal benchmark	125
74	Public API and portable reproduction	126
74.1	Exact types and immutable results	126
74.2	Package contents and replay commands	127
74.3	Adaptation and integrity boundary	128
75	Conclusions	128
76	Byte identities	128
V	Event-budgeted quartic certificates for the ordered rule (Report 20)	129
	<i>Report 20: Event budgeted quartic certificates for exact finite CA iteration</i>	129
77	Result and implementation boundary	129
78	Source data and inherited local semantics	131
78.1	Constants and costs	131
78.2	What one factor means	131
78.3	Arithmetic factor indices	131
79	Canonical quadratic primitive compiler	132
79.1	Exact primitive ledger	133
80	Fixed candidate slots without a factor array	133
80.1	Paying for source selection	134
81	Exact simultaneous evaluation of each candidate	135
81.1	Raw keys and isolation	135
81.2	Finite context guards	135
81.3	The simultaneous output and canonical order	135
82	Least future change and the completion rounds	136
83	Proof of the complete natural fiber statement	137
84	Fully charged size and expansion bounds	137
84.1	Independent per macro construction check	138
84.2	Source preprocessing is separate	139
84.3	The universal-source count is only a slot count	139
85	Coordinate heights and bit complexity	139
86	Complete implemented mass two example	140
86.1	A coefficient explicit artifact rather than an abbreviation	141
87	Reproducibility and audit evidence	141
87.1	Pinned dependencies and provenance	142
88	Limits and the next concrete improvement	143
VI	The parallel two-involution rule: replacement, sparse evaluation and quartic certificates (Reports 26, 27 and 28)	144

<i>Report 26: Two parallel conservative involutions</i>	144
89 Main result and the object that changes	145
90 Prospective isolation on the full shift	146
90.1 Write supports and read windows	146
90.2 Selection without a recursive eligibility test	146
91 The finite source and five particle geometry	148
91.1 Source interface and admissibility	148
91.2 The unsigned micrograph	148
92 Literal templates and the new two blocks	149
92.1 Edge templates	149
92.2 Phase templates	149
92.3 Raw predicates and guard reads	150
93 Preservation on every admissible state	150
93.1 Unique recognition	150
93.2 The complete corridor boundary partition	150
93.3 Every interaction and every guard	151
93.4 The phase block and agreement theorem	152
94 Clocks observations and time symmetry	152
94.1 Time reversal and the exact path	152
94.2 Halt observer and clean exact targets	153
95 Radius and finite resources	153
96 Malformed inputs and why prospectivity matters	154
97 Executable evidence and its limits	155
97.1 Reference interpreter and checks	155
97.2 Finite results	155
97.3 What transfers to arithmetic certificates	156
98 Relationship to prior work	156
99 Portable package and verification	157
100 Further questions	158
<i>Report 27: Sparse evaluation of a two involution particle cellular automaton</i>	160
101 Result and semantic boundary	160
102 Source data and endpoint geometry	161
102.1 Accepted sources and expanded metadata	161
102.2 Literal endpoint families	161
103 The new parallel rule	162
104 Complete sparse candidate discovery	163
104.1 Recovering every possible type	163
104.2 Recovering every raw key and orientation	164
105 Selection equality and local execution	165
106 Work space and bit bounds	166
106.1 A bound before prospective rejection	166
106.2 Expanded source and runtime costs	166

106.3 Integer lengths and deterministic bounds	167
107 Supported API and reproducibility contract	167
108 Measured evidence and independent checks	168
108.1 Small-source differential testing	168
108.2 Frozen universal-source benchmark	168
108.3 Frozen packet and release verification	169
109 Attribution and arithmetic certification limits	169
110 Offline replay and artifact provenance	170
<i>Report 28: A fixed horizon quartic compiler for the parallel particle rule</i>	172
111 The precise result	172
112 Finite source and endpoint geometry	173
112.1 Validated source data	173
112.2 The complete endpoint families	174
113 The new parallel rule and its inverse	175
114 Direct source slots and exact type identifiers	175
115 Sorting isolation and deterministic lane compaction	177
116 Prospective rediscovery and simultaneous transport	178
117 Quadratic arithmetic and complete witness uniqueness	179
117.1 Signed assignment	179
117.2 Exact integer comparison	179
117.3 Boolean operations and selections	180
117.4 Acceptance checks and the quartic	180
118 Resources that are actually charged	181
119 An explicit bound on every witness coordinate	182
120 Concrete emitted examples	183
120.1 A fully expanded two-particle quartic	183
120.2 Measured mass dependence at one whole step	184
120.3 The malformed old versus new separator	185
121 Boundary cases and interface semantics	185
122 What the finite computations establish	185
122.1 Baseline and nonvacuous exhaustive domains	185
122.2 Supplements and independent audit	186
122.3 Timing and verification limits	186
123 Release contents and independent replay	187
123.1 Requirements and read-only procedure	187
123.2 Scientific dependency boundary	188
124 Further research with distinct obligations	188
125 Conclusion	188
126 File map for the scientific argument	189

VII Smaller sufficient radii for the parallel rule: two-scale recognition and shorter exactness windows (Reports 70 and 71)	190
<i>Report 70: Two scale recognition for the binary five particle compiler</i>	190
127 Objects and principal result	191
128 The two scale full shift lemma	192
128.1 Recognition and reversible local control	192
128.2 Finite local meaning of the prospective test	193
128.3 Recognition invariance on arbitrary inputs	193
128.4 Every selection decision is invariant	194
128.5 Involution conservation and radius	194
129 A reversible choice among templates	194
130 The source and five particle geometry	195
130.1 The complete source interface	195
130.2 Signed modes and encoding	196
130.3 The precise admissible graph	196
131 All literal templates and local predicates	196
131.1 Edge templates	196
131.2 Raw recognition and guarded control	197
131.3 The unchanged phase block	198
132 Agreement on every admissible state	198
132.1 Every raw match uses the actual head	199
132.2 Homes and same anchor branch alternatives	199
132.3 Both signed corridor partitions	199
132.4 Interaction walls and old marker anchors	200
132.5 The two endpoints keep the same recognized anchor	201
133 Exact clocks observations and targets	201
134 The exact radius and resource ledger	202
134.1 The inherited universal numerical corollary	203
135 A four particle input separates the full shift rules	204
136 A separate complete radius two certificate	205
136.1 Conservation is an integer graph potential	205
136.2 A complete finite enumeration	206
136.3 All nonprojections have explicit collisions	206
136.4 What the classification does and does not imply	207
137 Dependencies certificate transport and attribution	207
137.1 What transports from earlier reports	207
137.2 Established architectural precedents	208
138 Evidence and reproducible presentation	208
139 Further mathematical directions	209
<i>Report 71: Shorter exactness windows for the binary five particle compiler</i>	211
140 Result and comparison	211
141 Source, geometry and complete endpoint families	212
141.1 Source interface and admissibility	212
141.2 Literal edge endpoints	213

141.3	Literal phase endpoints and exactness	213
142	Full-shift definition and safety of both new blocks	214
142.1	Edge: mutual guarded uniqueness and raw-anchor selection	214
142.2	Phase: typed keys and its own prospective filter	215
143	All-anchor admissible agreement	216
143.1	The close pair and the only possible singleton	216
143.2	Homes and multiple raw branch types	217
143.3	Both signed corridors and every equality case	217
143.4	Every interaction and the old endpoint key	218
143.5	The desired edge and phase selections all pass	218
144	Radius, unchanged dynamics and exact scope	219
144.1	A support refinement for the very same rule	219
144.2	Clocks, observers, targets and dependencies	220
145	Actual malformed inequalities for both blocks	220
145.1	Triple shortening changes each block separately	220
145.2	The smaller phase exclusion also changes behavior	221
146	Why the retained pair padding matters	221
147	Evidence, provenance and reproducible presentation	222
147.1	Historical preservation is qualified precisely	222
147.2	What was executed and what authentication establishes	223
148	Conclusion	223
149	Further questions and research for Part VII	224
A	Provenance of the merge	225

1 The series and this report

[write] Sections 1–3 and Appendix A were written when the six manuscripts were merged into this report (batch 82, cluster M1, 3 October 2026); text added inside the manuscripts’ own sections at that time is marked [write]. Parts V and VI (Reports 20 and 26–28; batch 82, cluster M2) were added by a second write on the same day, which extended these sections and the appendix and marks its additions to the manuscripts in the same way. Part VII (Reports 70 and 71; batch 91, cluster W4) was added by a third write on 4 October 2026, which extended these sections and the appendix and marks its additions in the same way. Everything else is the text of the twelve manuscripts, with labels prefixed and the corrections listed in Appendix A.

1.1 Twelve Reports on one automaton

The twelve manuscripts are “Research Reports” 14–20, 26–28, 70 and 71 of one AI-assisted research pipeline. They concern one object: the binary, number-conserving, globally reversible cellular automaton that Report 15’s compiler builds from a reversible two-counter source, together with the explicit universal source that Report 16 feeds it. Report 14 is a sibling construction in the wider class of binary conservative (not necessarily reversible) automata. Reports 20 and 26–28 certify that automaton, replace its ordered execution by a parallel rule with the same admissible dynamics, and evaluate and certify the new rule; Reports 70 and 71 replace the new rule by two further extensions of the same admissible dynamics with smaller sufficient radii.

Report	archive	delivered title (pages)	here	labels
15	15	<i>Globally reversible binary computation with five particles</i> (21)	Part I, first	fpa:
14	23	<i>Five particles in a binary conservative cellular automaton</i> (24)	Part I, second	fpa:bc:
16	13	<i>A literal universal reversible two counter source</i> (26)	Part II	fpa:ls:
17	16	<i>Startup history in a literal reversible source</i> (14)	Part III, first	fpa:st:
18	04	<i>Cellular clocks for a literal reversible source</i> (13)	Part III, second	fpa:cc:
19	09	<i>Exact lazy finite support evaluation of the frozen reversible binary CA</i> (14)	Part IV	fpa:lz:
20	08	<i>Event budgeted quartic certificates for exact finite CA iteration</i> (16)	Part V	fpa:eb:
26	20	<i>Two parallel conservative involutions</i> (16)	Part VI, first	fpa:pi:
27	18	<i>Sparse evaluation of a two involution particle cellular automaton</i> (11)	Part VI, second	fpa:sp:
28	03	<i>A fixed horizon quartic compiler for the parallel particle rule</i> (19)	Part VI, third	fpa:pq:
70	91:22	<i>Two scale recognition for the binary five particle compiler</i> (21)	Part VII, first	fpa:tr:
71	91:23	<i>Shorter exactness windows for the binary five particle compiler</i> (13)	Part VII, second	fpa:sw:

Report 15 is printed first because it is the base of the merge and because Reports 16–19 build on its compiler; Report 14 was written about forty minutes before Report 15 (archive member times 19:58 and 20:38 on 2 October 2026) and neither cites the other. Reports 20 and 26–28 follow in Report order: Report 20, the only one of the four about the old ordered rule, has a Part of its own, and the three Reports on the new rule share Part VI. Archive numbers are those of batch 82, except “91:22” and “91:23”, which are manuscripts 22 and 23 of batch 91. Reports 70 and 71, delivered a day later, follow in Report order as Part VII; the Reports numbered in between

concern other objects (Section 1.2). The archive names and the files each Report contributed are listed in Appendix A and in the report’s `README.md`.


1.2 Where the rest of the series is printed

The Reports cite one another by number, and so does this report. The whole series is filed in the Hilbert’s-tenth-problem category of the research-report collection (directory `SetTheory/Cardinals/docs/reports/hilbert-tenth-problem/`):

Reports	where printed
8–12	<code>signal-machine-collision-certificates</code> (SMC), sources 13–17: Report 8 is source 13 (Part IV), Reports 9–10 are sources 14–15 (Part V; Report 10 is the clean-target theorem), Reports 11–12 are sources 16–17 (Part VI; Report 12 is the four-mass decision theorem)
13	SMC Part VI, source 18 (the timed four-mass quartic; batch 82M4)
14–19	this report, Parts I–IV
20, 26, 27, 28	this report, Part V (Report 20, event-budgeted quartic certificates) and Part VI (Reports 26–28: two parallel involutions, sparse parallel evaluation, fixed-horizon parallel quartic certificates); batch 82M2
21–22	SMC Part VIII, sources 22–23 (the three-mass automaton; batch 82M2)
23–25, 33, 34	<code>fixed-universal-polynomials</code> (batch 82M3)
29–31	SMC Part VII, sources 19–21 (mass four in every dimension, orbit geometry, a binary planar shuttle; batch 82M4)
32	<code>group-theoretic-substrates</code> , Part V (a universal matrix semigroup; batch 82M3)
66, 67, 69	SMC, placed for its Part XIII (native-gap certificates and compilers; batch 91, <code>d750d98dd</code>); Report 70’s presentation tools adapt those of Reports 66 and 68
68	never delivered to ProveIt; only its <code>README</code> survives, as SMC’s <code>38-1ow-arity-qa-pred-tools-README68.md</code> (delivered inside the archive of Report 69)
70, 71	this report, Part VII (two-scale recognition and shorter exactness windows; batch 91)


Some of these destinations were placed in the same batch and may be written after this report; they are cited by path and Report number. The rows for Reports 66–71 were added by the third write (4 October 2026); Reports between 35 and 65 of the same pipeline are filed elsewhere in the collection or were not delivered, and none of them is cited by this report.

1.3 How the Reports depend on one another



An arrow points from a Report to one that uses it. Report 12 supplies the matching lower bound of the threshold corollaries of Reports 14 and 15; Report 14 compiles the 8,408-row two-

counter program of QOC Part V; Report 16 instantiates the universal source that Report 15 left existential; Report 17 changes four history tags of Report 16's source; Report 18 lifts Report 17's comparison to the automaton's clock; Report 19 evaluates Report 15's automaton for Report 17's source. Report 18 also reuses Report 16's shared-offset certificate. Inside the series, Reports 20 and 26–28 (Parts V and VI) use Reports 15, 17 and 19, and Report 28's bundled compiler proof also cites Report 14; SMC source 19 (Report 29) cites Report 16 for its four-versus-five boundary; GTS Part V (Report 32) uses Report 16's loader and Turing table. The four Reports of Parts V–VI depend on one another and on the earlier Parts as follows.



Report 20 certifies Report 15's ordered rule and uses Report 19's candidate lemma; Report 26 keeps Report 15's templates, replaces their execution, and takes Report 19's universal ledger; Report 27 evaluates Report 26's rule; Report 28 certifies it and uses Report 27's evaluator as a test oracle. Reports 70 and 71 (Part VII) keep Report 26's templates and admissible graph, compare their bounds with Report 26's, and use Report 19's universal ledger; Report 70 keeps Report 26's phase block, and Report 71 pins Report 70's release and replaces both blocks. The dashed box is the arithmetic layer that Reports 20 and 28 prove again: Report 20 cites it as Report 8 (SMC source 13), Report 28 cites only Knuth, and CDC is credited in notes.

1.4 What the Parts prove

- **Part I (Reports 15 and 14).** Every finite partial-injective guarded two-counter machine with a terminal halt compiles to a globally reversible (bijective on $\{0,1\}^{\mathbb{Z}}$), number-conserving binary automaton of explicit finite radius, with inputs of exactly five ones and an anchored halt word (Theorem 4.1); with Report 12 this makes five the exact particle threshold for globally reversible binary automata, for halt-pattern occurrence and for exact finite targets (Corollary 16.1). Report 14 proves the same number five for binary conservative automata by a different compiler for every deterministic ADD/SUB/NOP program (Theorem 18.1, Corollary 27.1), with a literal universal instance of radius 756,787 and a persistent halt word. Both give fixed-horizon degree-four sum-of-squares certificates with empty-or-singleton natural fibers and paid nonnegative-real variants.
- **Part II (Report 16).** One pinned reversible two-counter machine with 122,622 controls and 141,561 literal rows is universal on a clean loader family (Theorem 33.1); on that family positive return and periodicity of the fixed five-particle automaton are r.e.-complete, with least period exactly $2\Theta + 2$ (Corollary 38.1).
- **Part III (Reports 17 and 18).** Four history-tag swaps give a new source with the same semantics whose literal clock is no larger at every normalized boundary and strictly smaller at every Turing-machine cut (Theorem 42.1); first encounters determine the optimal static orientations (Theorem 47.1); both statements hold for the automaton's clock as well (Theorems 53.1 and 53.2), through the identity $\Theta = \lambda M + \text{TV}(A^2) + \text{TV}(B^2) + U$.

- **Part IV (Report 19).** A source-sized lazy evaluator equals the frozen ordered product of 269,291,358,255 involutions on every finite support, inverse included (Theorem 65.1).
- **Part V (Report 20).** For a fixed accepted source, mass n , horizon T and budget K , a sum of squares of quadratic residuals has a singleton natural auxiliary fiber exactly when the literal ordered product completes T steps with at most K changing factors (and meets a prescribed endpoint, if any), and an empty one otherwise (Theorem 77.1). It uses $T + K$ rounds over $C = \binom{n}{2}[\Delta + 3 + 4\max(n - 2, 0)]$ occupied-pair candidate slots (170 at mass five for the universal source) and no array over the F factors. Implemented are a scheduler reference and a mass-two exporter (612 witnesses, 616 residuals, 4,573 expanded monomials); an arbitrary-mass exporter is not.
- **Part VI (Reports 26–28).** Report 26: $F_{\text{new}} = P_{\parallel} \circ E_{\parallel}$, a product of two full-shift, number-conserving involutions, agrees with the ordered rule on the entire admissible doubled micrograph, has sufficient radius $180D + 258 + 9J$ (91,711,698 for the universal source) and is time-symmetric (Theorem 89.1), by a prospective-isolation lemma of radius $3(b + r)$ (Lemma 90.1); it differs from the ordered rule on malformed inputs. Report 27: a sparse evaluator equals the new rule on every finite support, with at most $n + 2$ raw-discovery calls per step (Theorem 101.1), and runs 128 universal startup steps. Report 28: for a fixed source, mass n , horizon T and direction, a polynomial of degree at most four in $4n$ external naturals has exactly one natural witness at accepted endpoint pairs of the new rule and none elsewhere (Theorem 111.1); its witness count $4\max(n - 1, 0) + T(w_E + w_P)$ grows with T , and its two-particle example has 1,494 witnesses, 1,502 residuals and 12,595 monomials.
- **Part VII (Reports 70 and 71).** Report 70: a full-shift involution lemma with separate write, recognition and control radii $b \leq a \leq r$, isolation at $H = \max\{2(b + a), b + r\}$ and radius at most $b + a + H$ (Theorem 128.1); with Report 26's phase block it gives $F_* = P_{26} \circ E_*$, which agrees with Report 26's rule on the whole admissible doubled micrograph and has sufficient radius $108D + 149 + 3J$ (55,027,013 for the universal ledger; Theorem 127.1); an explicit four-particle input separates F_* from Report 26's rule; independently, the injective number-conserving binary rules of radius two are exactly the five shifts (Theorem 136.1). Report 71: with triple exactness shortened from $3b + 1$ to b in both blocks, $F_s = P_s \circ E_s$ agrees with both predecessors on admissible states and has sufficient radius $76D + 105 + 3J$, or $76D + 104 + 3J$ for the same rule (38,722,713 and 38,722,712; Theorem 140.1); malformed inputs separate each new block from both predecessors, and the pair exactness L cannot be changed for these templates and this selector (Section 146).

1.5 Questions answered inside the series and elsewhere

Inside this report: Report 15's first research question (a literal universal reversible source) is answered by Part II; Report 15's limit on periodicity conclusions (the missing no-blocking premise) is supplied for Report 16's clean family by Corollary 38.1; Report 16's first two questions (initialization cost; a lazy indexed compiler) are answered by Part III and Part IV; Report 17's open old/new cellular-clock comparison is settled by Report 18, and its lazy-evaluator question by Report 19. In Parts V–VI: Report 26's second question (a lazy evaluator for the new rule) is answered by Report 27, its third (an arithmetic certificate for the new selection mechanism) by Report 28 at each fixed horizon, and Report 27's arithmetic-verifier question by Report 28; Report 20's next step (an arbitrary-mass emitter for the ordered rule) and Report 16's certificate-tradeoff question stay open. In Part VII (4 October 2026): Report 26's first question, whether type-specific support geometry can reduce $H = 2(b + r)$, is answered affirmatively for the compiler's templates by Reports 70 and 71, each for a new full-shift rule; Report 70's further

directions (a smaller common exclusion bound; a practical evaluator or arithmetic certificate for its rule) stay open. Dated notes stand at each of these questions.

Reports 20, 26–28, 70 and 71 answer no question of a neighbouring report. Reports 14–19 bear on questions of two neighbouring reports (labels of those reports are given in typewriter type):

question	Report	effect
SMC Part IV, § <code>smc:sl:sec:questions</code> , item 1: “Print one safe universal source”	16	answered in substance for Report 15’s guarded syntax (122,622 controls, 141,561 rows, clean loader, all-input proofs); the lattice table of SMC Part IV in source 13’s pulse interface is still not printed
SMC <code>smc:tm:q:source</code> : a universal reversible two-counter table in the separated syntax	16	answered in part: a literal universal reversible two-counter table now exists, but not in source 14’s separated syntax
SMC: “a reversible numerical threshold, a binary upper bound” (front matter; Part IV item 4); source 17’s non-claim “a five-particle upper bound for binary automata”; the mass-5 cell of SMC’s threshold table	14, 15	a binary upper bound five (Reports 14 and 15); with Report 12, five is exactly the threshold for binary and for globally reversible binary number-conserving automata (Corollaries 27.1 and 16.1); source 17’s non-claim remains true of source 17
SMC <code>smc:tm:q:mass</code> , <code>smc:ct:q:mass</code> : the sharp reversible threshold	15	answered for automata with one unit symbol (binary); with several weight-one types three suffices (SMC Part V)
QOC <code>qoc:um:thm:poly</code> : certificates for the 8,408-row source	14	bears on it: Report 14’s certificate for the same source is weaker (degree four, not two) but adds the automaton’s clock (note at Section 28)

1.6 Relation to the collection and to the formal project

Re-proofs. Several passages re-derive results already printed in the collection; they are kept, as second presentations, with a [write] pointer naming the host statement and without any novelty claim: Report 14’s Section 26 (the 528- and 8,408-instruction programs and their proofs are QOC Part V’s Lemmas `qoc:um:lem:virtual` and `qoc:um:lem:prime`); the Neary–Woods table and its (u_{10}, b) discrepancy in Reports 14 and 16 (printed once in QOC Part III, after Table `qoc:wf:tab:tm`); Report 16’s front end (Section 34); Report 15’s clean-target wrapper (Section 13), which is the forward/inverse/bridge construction of SMC source 15 (Report 10, Theorem `smc:ct:thm:restore`) for guarded sources, and is Bennett’s uncomputation [8]; Report 16’s history and prime-encoding invariants (Sections 35–37), which make literal the existence argument of SMC Part IV, §`smc:sl:sec:safe`. Within this report the compiler clock $\tau_e(c) = 3 + 2(Z + c) + \Delta_e - 4S$ is printed in five Reports (equation (5), (29), Report 17’s Section 50, (54.1), and Report 26’s (94.1)), and Report 19 restates Report 15’s ledger; the later occurrences carry pointers. In Parts V–VI: the circuit-to-quartic layer of Reports 20 and 28 (signed pairs, comparison, equality, Boolean gates, topological uniqueness) is the canonical circuit lemma `cdc:pt:lem:circuit` of `canonical-diophantine-certificates` (CDC) with `cdc:pt:lem:signed` and `cdc:mem:lem:compare`, printed also in SMC Part IV (`smc:sl:eq:cmp`, `smc:sl:eq:eq`, and Report 20’s constant division in §`smc:sl:sec:presburger-quartic`); Report 28’s bitonic sorter is CDC’s Proposition `cdc:mem:prop:comparator-cost` and §`cdc:mem:subsec:bitonic`; Report 27’s Lemma 103.1 restates Report 26’s Lemma 90.1, and Reports 27 and 28 restate Report 26’s rule and radii. Neither Report 20 nor Report 28 cites CDC; each of these passages carries a [write] credit, and none is claimed as new. In Part VII: Report 70’s Lemma 131.1 restates Report 26’s phase block; its Sections 130, 131 and 133 restate the source interface, templates, clocks and clean targets of Reports 15 and 26, and Report 71’s Section 141 restates

Report 70’s with its new exactness predicate; Report 70’s conservation criterion (Lemma 136.2) is the radius-two analogue of the de Bruijn-graph potential that SMC source 16 (Report 11) uses for the elementary rules (§smc:su:sec:checks), and its count of 428 conservative tables is Fuk  and Sullivan’s [34]. Each of these passages carries a [write] pointer.

Reviews. At the time of the first write (cluster M1) the Hilbert’s-tenth research programme (Computability/HilbertTenthProblem/) had *not* reviewed Reports 14–19; for Report 16 see the dated note at the end of this paragraph. Its note `Papers/research-wip/native-stream-queue/reviewed_report_archive_replay.md` (commit ceb7222e9), which restores the batch’s archives from an earlier revision for replays, lists the six archives of this report as “Relocation only” and states that relocation is not a scientific review. Its review `Papers/research-wip/native-stream-queue/review_parallel_particle_reports.md` (commit fb7e3cb47) covers Reports 26–28 (Part VI), found no defect, and notes that the inherited template geometry and source simulation of Reports 15–17 remain a dependency it did not reconstruct; its review `review_binary_planar_four_particle31.md` (commit 16f50dcc6) concerns Report 31 (SMC Part VII) and separates it from the five-particle line. The programme’s note `neary_woods_explicit_universal_tm.md` transcribes the same Neary–Woods table. [write] (3 October 2026, second write.) Since the first write, the programme’s `review_literal_universal_reversible_source16.md` (commit 69730d3e8) has reviewed Report 16: it scans all 141,561 primitive rows of the pinned source independently, finds no source-interface or arithmetic contradiction, and records Report 16’s literal-source and clean-nonblocking obligations as closed, while keeping Report 15’s compiler theorem and the Neary–Woods universality theorem as explicit dependencies. Three scoped intakes followed, each reading the proofs without executing the archived Python: `review_cellular_clock_domination18_intake.md` (commit 735f62a6a) finds that Report 18’s central clock-domination proof passes and closes the CA-clock comparison for the fixed history-bit orientation class; `review_lazy_ordered_evaluator19_intake.md` (commit c875bad40) finds no gap in Report 19’s finite-support argument, as a proof intake rather than an implementation audit; and `review_event_budget_clean_clock_intake.md` (commit 8aabb1453) reads Report 20’s complete proof (Part V) and identifies no new defect within its stated scope. All three stress what the Reports do not supply: a smaller arithmetic compiler, a paid ordinary-input loader, or an unbounded fixed-arity history representation. Reports 14, 15 and 17 have no review. [write] (4 October 2026, third write.) The programme’s scoped intake `review_new_arithmetic_0d7f51c44.md` (commit a9ab9a698) read Reports 70 and 71 in part, found no correction, and states that Report 70’s full-shift theorem and Report 71’s recognition and geometry proof were not re-audited (Part VII).

Formal status. No statement of this report is formalized in Lean or Rocq, and none of the Reports cites a declaration; the report’s placement in the collection confers no formal status. Report 15’s research question 5 and Report 16’s question 4 propose mechanizing the trusted cores; Report 14’s question 4 proposes formalizing its full-shift extension; Report 26’s question 4 proposes formalizing its lemma and corridor partition, and Report 70’s last section its anchor and signed-corridor arguments.

2 Conventions and notation

Labels and numbering. Every label carries the prefix `fpa::`. Report 15’s labels are `fpa:` followed by the delivered name, and the other Reports use the sub-prefixes `fpa:bc:` (14), `fpa:ls:` (16), `fpa:st:` (17), `fpa:cc:` (18) and `fpa:lz:` (19); labels added in the write are those beginning `fpa:sec:`, `fpa:part:` and `fpa:app:`, the labels `fpa:rep:r:first` and `fpa:rep:r:last` on the first and last section of Report *r*, and `fpa:lz:sec:pinned` and `fpa:lz:sec:bytes` on two unlabelled sections of Report 19. Parts V–VI use the sub-prefixes `fpa:eb:` (Report 20), `fpa:pi:` (26), `fpa:sp:` (27) and `fpa:pq:` (28) followed by the delivered

names; their write added `fpa:part:five`, `fpa:part:six`, the labels `fpa:rep:r:first` and `fpa:rep:r:last` for $r = 20, 26, 27, 28$, and `fpa:eb:sec:source`, `fpa:eb:sec:proof` and `fpa:pq:sec:sorting` on three unlabelled sections. Part VII uses `fpa:tr:` (Report 70) and `fpa:sw:` (Report 71) followed by the delivered names. The two Reports deliver eleven identical label names (`thm:main`, `lem:phase`, `eq:ahead`, `eq:behind`, `eq:commit`, `eq:constants`, `eq:direct`, `eq:dispatch`, `eq:endpoint`, `eq:free`, `eq:mutual`), which the two sub-prefixes keep apart. The third write added `fpa:part:seven`, `fpa:rep:r:first`, `fpa:rep:r:last` for $r = 70, 71$, and `fpa:sec:partseven:questions` on its closing section of further questions. Sections and theorems are numbered continuously through the report (Theorem $n.m$ is the m th numbered statement of Section n). Reports 15 and 16 number their displays by hand, (1), (2), ..., and their prose cites them as “(3)”, “(26)”; these numbers are kept as delivered and are *local to the Report*: in Part I a bare (n) is Report 15’s display n , in Part II it is Report 16’s. Displays that the Reports number automatically (Reports 14, 17, 18, 20, 26–28, 70 and 71) are numbered $(n.m)$ by section, so they cannot be confused with the hand-numbered ones.

Delivery paths. The Reports name their files by their paths inside the delivered archives (for instance `source/source.json`, `reproducibility/run_checks.py`, `companions/report15/`, ...). Those names are kept in the text. The `README.md` of this report maps every delivery path to its shipped name (shipped files carry the prefixes `14-five-binary-`, `15-rev-five-`, `16-literal-source-`, `17-startup-`, `18-cell-clock-`, `19-lazy-eval-`, `20-event-budget-`, `26-parallel-involutions-`, `27-sparse-parallel-`, `28-parallel-quartic-`, `70-two-scale-` and `71-short-windows-`), and lists the files that are not shipped: byte copies of other files, the delivered PDFs and checksum ledgers, the companion copies of Reports 12 and 15, and eighteen heavy generated files (among them both `source.json` tables, each 32,034,272 bytes) that the shipped programs rebuild in seconds (README, “Reconstructing the excluded data”). Parts V–VI exclude in the same way Report 27’s copy of the universal source with its two universal traces (and their byte-identical optimized copies), and Report 20’s gzip of the source. Part VII does not ship the page renders of the two delivered PDFs, the checksum ledgers, the container ZIPs whose every member is shipped, or the byte copies of Report 26’s archive and of files already shipped by Parts II–VI (README).

Notation. No symbol of a Report was renamed: each Part uses its Report’s symbols. The same letter often means different things in different Reports; the table lists the collisions. In the text written for the merge, these notions are named in words.

symbol	Reports 15–19 (Parts I–IV)	Report 14 (second half of Part I) and others
m	number of source controls	same (all controls, halt included)
p	number of nonzero-update branches	number of ADD or SUB <i>rows</i> (a SUB row has two branches)
a	number of zero-update branches	a counter value ($c_0 = a$, $c_1 = b$; initial a_0, b_0)
n, z	Report 19: n is the mass of a support	n NOP rows, z SUB rows
D	$2m + 4p$, the gap alphabet of signed modes	$m + 2p$, the gap alphabet of unsigned codes
S	$2D + 2$, the home offset of the close pair	—
Z	$10B_3 + 10 + 2J = 40D + 60 + 2J$, the counter buffer	$4K = 16D + 20$
K	Report 16: the source horizon in primitive rows of the certificates (141,565 K witnesses)	$4D + 5$, the component and contact distance
L, R	Report 15: travel constant $L = 3D + 4$ and radius bound R , eq. (18); Reports 16–18: half-tape integers L, R	radius $R = 30D + 37$; L, R also half tapes in Section 26
T	Report 15: context radius of a gate, and a microedge count in Proposition 10.1; Reports 16–18: the scratch counter	scratch register; tick and clock variable in the certificates

symbol	Reports 15–19 (Parts I–IV)	Report 14 (second half of Part I) and others
H, h	Report 15: source horizon H of a certificate and halt control h ; Reports 16–18: history counter H ; H a home mode	horizon h of a certificate; H the terminal residual
Θ	the automaton’s clock: microedges to the first plus-halt home	—
λ	Report 18: $3 + 2Z - 4S = 36,684,691$ for the universal source	—
J	the class cutoff of guards	a Turing state ($J1 = (u_{10}, b)$ is the halt cell), as also in Report 16’s Table 16
A, B	Reports 16–18: the two physical counters; Report 15: A_t a residual, B a swap radius or normalized-branch count	B the number of branches
F	Report 15: one automaton step $P \circ E$; Report 19: the number of factors	—
M, U	Report 18: moving-row and zero-update counts; Report 15 certificates: M bounds the initial counters	U_t the one-hot residual

In Parts V–VI the constants $m, p, a, D = 2m + 4p, S = 2D + 2, L = 3D + 4, Z, J$ and $F = 8pD + 29p + m + a$ mean what they mean in Reports 15–19 (Report 20 writes S_0, L_0 for S, L), and n is always the mass. The second table lists the collisions of these four Reports, among themselves and with Parts I–IV. One printed symbol changed: Report 20’s height bound H_* is printed Ht_* .

symbol	Report 20 (Part V)	Reports 26–28 (Part VI)
$\mathcal{F}$	the <i>old ordered</i> rule, one CA step	Report 28: the <i>new parallel</i> rule (Report 26 writes F_{new} , Report 27 Φ ; Report 19 writes $\mathcal{F}$ for the ordered rule)
F	the number of ordered factors	the number of endpoint types (the same number $8pD + 29p + m + a$; Report 26: the factors become alternatives inside two blocks)
b	the number of branches $p + a$ (Report 27 writes t)	$b = B_3 = 4D + 5$, the triple write radius (Report 15’s B_3)
H	$H = h - z$, a signed head-to-marker offset; Ht_* (delivered H_*) the witness-height bound; H_0 the largest input coordinate	$H = 2(b + r)$, the isolation distance (H_E, H_P); Report 28: $H_* = \max(H_E, H_P)$; H_q is a home mode throughout
C	the number of candidate slots	Report 27: a bound on discovered identities, $\mathcal{C}(X)$ the raw-key set; Report 28: $\mathcal{C}_B(X)$ the raw-key set, $C(s, n)$ the witness-height constant
K	the budget of changing factors	Report 26: $K = 3b + 1$, the triple exactness radius; Report 27: a bound on raw keys; Report 28: the slot widths K_E, K_P (Report 16: a horizon; Report 14: a contact distance)
T	the horizon	Report 28: the horizon; Report 26: the block map $T(x)$ of the lemma, and a path length; Report 27: the triple shape $T(M, \sigma, h)$ and a run time T_{raw}
r, R	$r = D + 3$, an index constant; $R = T + K$ the number of rounds	r the read radius (r_E, r_P); R a radius (R_{new}, R_E, R_P); Report 28: R_j the residuals
W	bits per auxiliary	Report 28: the witness count $W(T)$; Report 26: W_g a write set

symbol	Report 20 (Part V)	Reports 26–28 (Part VI)
E, P	E the number of residuals, P the polynomial; also the E and P factor blocks	the two involution blocks; Report 28: $P_{s,n,T,\epsilon}$ the polynomial, P_E, P_P padded widths
A, I, G, Q	ledger counts: assignments, comparisons, Boolean gates, acceptance rows	Report 28: the same ledger counts; Report 26: $A(x)$ the selected-key set; Report 27: $A(M, \sigma, h)$ a pair shape, A a coordinate bound, $Q = 1 + n^3 + n^2 \Delta$
Δ	the largest home incidence list	Report 27: the same; Report 26: the update $\Delta \in \{-1, 0, 1\}$ of a branch (Report 15's Δ_e)
λ, ℓ	λ the largest coefficient bit length	Report 27: λ an orientation, ℓ_g an exactness radius; Report 28: λ_g an exactness radius, λ a lane, ℓ both an orientation and $\lfloor n/2 \rfloor$ (Report 18: $\lambda = 36,684,691$)
ϵ	a parity correction	Report 27: the same; Report 28: the time direction ± 1
M, U	U a gate bound	Report 28: M the number of residuals, U a magnitude in $\text{RawCap}(U)$; Report 27: M the metadata size, U a loose bound (Report 18: M moving rows)
$\mathbb{N}$	$\mathbb{N}$, with zero	Report 26: $\mathbb{Z}_{\geq 0}$; Report 28: $\mathbb{N}_0$ (both with zero)

Tempting false readings in Parts V–VI: Report 20's $\mathcal{F}$ (the ordered rule) is not Report 28's $\mathcal{F}$ (the parallel rule), and the two rules differ on malformed supports; Report 20's b counts branches, while b in Reports 26–28 is a radius; Report 20's C counts slots, not Report 28's height constant $C(s, n)$; Report 26's $K = 3b + 1$ is a radius, not Report 20's budget.

In Part VII the constants $m, p, D = 2m + 4p, S = 2D + 2, B_2 = D + 1, L = 3D + 4, b = B_3 = 4D + 5, K = 3b + 1, Z = 10b + 10 + 2J$ and J mean what they mean in Report 26, and $\mathbb{N}$ is $\mathbb{Z}_{\geq 0}$. The third table lists the collisions of Reports 70 and 71 with the rest of the report. No printed symbol of either Report changed.

symbol	Reports 70–71 (Part VII)	elsewhere in this report
a, α	a the recognition radius ($3b+1$ in Report 70's edge block, b in Report 71's); α the number of zero-update branches	a the number of zero-update branches (Reports 15–28); Report 14: a counter value
r, r_P	$r = Z+J$ the control radius of the edge block; r_P the phase read radius (K in Report 70, b in Report 71)	Report 26: the read radius ($r_E = Z + J, r_P = K$); Report 20: $r = D + 3$
H, H_P	$H = \max\{2(b+a), b+r\}$, which is $b+r$ for the compiler; H_P the phase isolation distance, $2(b+K)$ in Report 70 and $4b$ in Report 71	Report 26: $H = 2(b+r)$; Report 20: a head-to-marker offset; H_q a home mode throughout
F_{26}, E_{26}, P_{26}	Report 26's rule and blocks, that is $F_{\text{new}}, E_{\parallel}, P_{\parallel}$ of Part VI; F_{70}, E_{70}, P_{70} (Report 71) are Report 70's F_*, E_* and P_{26}	F_{old} in Part VI is the <i>ordered</i> rule of Report 15, not F_{26}
E_s, P_s, F_s	Report 71's blocks and rule; s stands for “short”	$s \in \{-1, +1\}$ the side of a phase triple, also in Part VII; $s = \max\{r, b + 2a\}$ a read radius in Report 70
T_u	the local controlled involution at anchor u	T a horizon (Reports 20, 28), Report 26's block map $T(x)$
$\mathcal{R}, \mathcal{S}, \rho$	recognized-anchor set, selected-anchor set, recognition predicate (ρ_g for type g)	R a radius throughout; ρ a parity bit in Report 16

symbol	Reports 70–71 (Part VII)	elsewhere in this report
$C_u(x)$, $C_P(x)$ m, p	the guarded edge types at anchor u ; the typed phase-key set m the number of controls, but also the selected old marker m (Report 70) and m_0, m_1 (Report 71); p the nonzero-update branches, but an integer potential on de Bruijn vertices in Lemma 136.2	Report 26: $C(x)$ the raw-key set; Report 20: C the slot count m, p the counts only
f, g	Report 70’s radius-two appendix: f a five-input local rule and $g = f - c$ an edge weight; elsewhere g an endpoint type	g an endpoint type (Reports 26–28)
A, C, N, t, w	a corridor’s departure and arrival markers, its length $N = w(C - A)$, the oriented head position t , the direction w	Report 20: ledger counts A , slot count C ; Report 28: $W(T)$ the witness count
ℓ_g	the exactness radius of type g : $3B_g + 1$ in Report 70; L for pairs and b for triples in Report 71	Report 27: an exactness radius; Report 28: λ_g

Tempting false readings in Part VII: F_{26} is Report 26’s *parallel* rule F_{new} , not the ordered rule F_{old} of Part VI; a is a radius, not the zero-update count, which these two Reports call α ; the four bounds $180D + 258 + 9J$, $108D + 149 + 3J$, $76D + 105 + 3J$ and $76D + 104 + 3J$ are sufficient radii of three different full-shift rules (F_{26} , F_* and F_s , the last two bounds for the same rule) that agree on admissible states, not successive radii of one automaton; and they are spatial radii, not arithmetic gate counts.

Tempting false readings in Parts I–IV: Report 14’s $D = m + 2p$ is not Report 15’s $D = 2m + 4p$ (they count different alphabets, and p is not the same count); the two radius bounds $30D + 37$ (Report 14) and eq. (18) (Report 15, quadratic in D) bound different automata; Report 16’s K is a horizon, not Report 14’s contact distance; J in “J1” is a Turing state, not a class cutoff.

3 Status and non-claims

The twelve Reports are AI-assisted research manuscripts; Report 14’s delivery README says “prepared with AI assistance”. They are unrefereed and not formalized; the research programme has reviewed Reports 16, 18, 19, 20 and 26–28, by reading their proofs and running only its own checkers, has read Reports 70 and 71 in a scoped intake that did not re-audit their main proofs, and has not reviewed Reports 14, 15 and 17 (Section 1.6). Each Report states its own limits; all of them are kept in its text, and they are collected here.

- **Report 15.** Mathematical proofs with independent computational and mathematical audits, not proof-assistant formalizations or claims of novelty or priority. Universality is in an effective finite-input halting-observation sense; there is no claim of intrinsic universality, real-time simulation, efficient input conversion, radius optimality, or a printed universal binary truth table. The universal source of Corollary 4.2 is existential in Report 15 (Part II materializes one); the executable example is finite and accepting, not universal. The radius bound is an upper bound, and writing the truth table would be impractical. The certificate horizon is a compiler argument; no unbounded single-fold representation follows; the paid variant concerns the nonnegative orthant with natural initial counters, not unrestricted reals. Periodicity and positive return need an extra no-blocking premise for general sources. The threshold depends on Report 12’s separate proof. The Alhazov–Imai construction was

not inspected in full; an incomplete search is not a novelty argument; no first-construction claim is made.

- **Report 14.** No reversibility, intrinsic universality, efficient input preparation, efficient CA time or minimal radius is claimed. The universal truth table ($2^{1,513,575}$ rows) and the 1,272,679,260 indexed cases are specified, not materialized; universality rests on the Neary–Woods source and all-input proofs, not on table size. The horizon is fixed at emission; there is no unbounded single-fold representation and no universal-polynomial record; the $h = 1$ universal schema is not an accepting universal computation. The sharp threshold depends on Report 12, excludes infinite backgrounds, active zero-weight states, several unit labels, signed weights and nonuniform rules, and is not a priority claim; the literature search was limited. The loader is not a polynomial-time map in the length of the Turing-machine input. Finite tests do not replace the proofs.
- **Report 16.** An ordinary mathematical proof with executable audits, not a proof-assistant formalization, efficient simulation, novelty claim or size record. The automaton’s factor list, truth table and universal-instance quartic are not materialized. Empty-tape initialization has a predicted clock near $8 \cdot 10^{13}$ source steps; only small literal macros and the fresh entry were executed. Universality and no-blocking hold under the clean-input promise; the clean wrapper and its automaton are theorem-level only; the four-mass lower bound is imported.
- **Report 17.** No raw cellular-automaton speed comparison and no completion of a universal program is claimed (Report 18 adds the clock comparison); the old 79,936,151,060,302-step startup is derived, not executed; an equal-history internal state can be slower (28 versus 104 steps). Optimality is within the fixed templates. Report 16’s machine and artifacts are not revised; equal counts do not identify the two automata. Not a formalization, novelty, efficiency or size claim.
- **Report 18.** The comparison holds within the fixed static-orientation, recorder, prime-template and compiler class. No novelty, optimal universal-machine size, unrestricted speedup or host-runtime claim; no uniform ratio of arrival times, no equal-time alignment; no automaton orbit or giant rule array was executed; symbolic factor and radius counts are not allocated objects.
- **Report 19.** An implementation-equivalence result, not a new universality proof, a proof-assistant formalization or a novelty claim. No constant-time or polylogarithmic cost guarantee; the 128 universal microsteps do not complete startup or a Turing transition; no eager universal comparison was made; the benchmark’s resource figures are environment-specific.
- **Report 20.** A fixed-parameter finite-schema theorem, with no real-witness, unknown-horizon, finite-fold MRDP, mass-two universality, novelty or optimality claim. Only a scheduler reference and a mass-two coefficient exporter are implemented; the arbitrary-mass exporter is not. The prefactor 10000 and the height bound are conservative theorem-level bounds, not measured counts. The universal count $C = 170$ is a slot count, not witnesses, equations or operations, and no practical universal quartic exporter is claimed. No polynomial-in-log J preprocessing; no unconditional small budget on malformed supports, where cascades occur; uniqueness is over natural auxiliaries only; finite and mutation tests corroborate and do not prove. The macro allocation is not a claim about exact gate counts per macro; the height bound concerns accepting fibers, not adversarial submitted values; the construction does not inherit the lazy evaluator’s smaller test count; the release manifest is an integrity record, not a proof of authorship or authenticity.
- **Report 26.** Neither radius optimality nor novelty of the two-involution architecture is claimed; no universal new-rule evaluation or truth table is produced. Agreement holds on

admissible states only: the new rule differs on malformed inputs, so the radius comparison is between sufficient bounds of two different full-shift automata, and the new bound is worse for tiny sources (978 against 540). [write] (4 October 2026.) For that source Reports 70 and 71 give 581 and 409 (408 by Report 71’s same-rule refinement), the latter below 540; each bound belongs to its own new rule (Part VII). Universal-source statements are inherited conditionally; the universal-source theorem is not re-proved and no universal template array is materialized or executed. Report 19’s evaluator and Report 20’s certificates do not transfer to the new rule; Report 26 supplies no arithmetic compiler, fibers or gate counts. Finite tests are corroboration, the release seal is not a signature, and nothing is formalized.

- **Report 27.** An implementation theorem on finite supports. The universal benchmark does not finish startup, a Turing-machine step or a halting computation, and no eager universal oracle is built. No worst-case sublinear-in- F bound and no polynomial-in-log J claim; hash costs are expected, not worst case. No new Diophantine compiler, universal equation or gate count. The architecture is attributed to prior work; no priority, radius minimality or novelty. The universal source’s universality is inherited, not proved again; arbitrary-support equality rests on the proof, not on the malformed roundtrips; the timings are machine-specific observations, not guarantees. Full replay requires Linux; not a proof-assistant formalization.
- **Report 28.** A fixed-horizon family, not a fixed-arity unbounded-time representation, and no finite-fold MRDP consequence; no uniqueness over $\mathbb{R}$, $\mathbb{Q}$ or signed auxiliaries. It certifies an exact requested endpoint, not an observer or halting predicate; the witness-height bound covers accepted endpoint pairs only. No universal-source polynomial was emitted; only the pair example is fully expanded; the mass-five certificate tests at horizons two and three were not completed. Costs are current implementation costs, not lower bounds or improvements over certificates with other semantics. No polynomial-in-log J claim; the expanded verifier checks only the delivered polynomial; the code is research software, not formally verified; historical ordered-rule tests certify no new-rule support.
- **Report 70.** Sufficient bounds, not measured or minimum radii; no external novelty or priority, minimum universal radius, radius-three construction or universal implementation is claimed, and the marker, landscape and two-involution architecture is attributed to prior work. The universal source table is absent from the package; universality and normalization are inherited premises, neither re-proved nor executed; the universal number is an inherited-ledger substitution. The new full-shift rule differs from Report 26’s off the admissible set; old arbitrary-input evaluators, Diophantine certificates, ordered-factor budgets and circuit costs do not transfer, and no new-rule simulator, template array, truth table or arithmetic evaluator is supplied; effective existence of the local rule is not a practical evaluator or a template-independent running-time bound. The radius-two appendix is independent of the main proof and limited to uniform, time-independent, symmetric radius-two binary rules (no claim for larger neighbourhoods or for partitioned, time-dependent, second-order, nonuniform or larger-alphabet models); the count 428 is established literature; the Schranko–de Oliveira full text was not obtained and no priority over it is claimed. Not a proof-assistant formalization; the release manifest is an internal consistency check, not a signature, and reproducible typesetting is not a proof.
- **Report 71.** Sufficient bounds, not measured or minimal radii; the one-site support refinement changes no rule, and the phase boundary predicates still read the full old write interval. The universal source table is absent; universality and normalization are inherited premises. The pair-padding counterexamples concern only the unchanged templates and selector with an enabled nonzero branch, not other encodings or rules, and are not failures of involutivity. No new-rule evaluator, arbitrary-input arithmetic certificate, priority claim or proof-assistant formalization; it corrects neither predecessor. Its source packet observed

concurrent changes in the live Report 70 tree; preservation records are endpoint observations, not atomic snapshots.

Part I

Five particles suffice in binary automata (Reports 15 and 14)

Globally reversible binary computation with five particles

An explicit finite-radius compiler and counted finite-horizon Diophantine certificates

Report 15 Mathematical construction and reproducible audit

3 October 2026

Abstract of Report 15, as delivered. A finite partial-injective deterministic machine with two natural counters is compiled into one time-independent cellular automaton on the full binary shift. Every encoded configuration has exactly five occupied sites and an otherwise zero background. The automaton is globally reversible, conserves the ordinary number of ones, and has an explicit finite radius. Full-shift reversibility is proved by an isolated equal-mass pattern-swap lemma valid on arbitrary dense or infinite inputs; simulation correctness is a separate matching argument on a precisely delimited admissible set. Forward and backward copies of a subdivided counter computation reflect at missing edges. The classical existence of a universal reversible two-counter source gives a five-particle universality existence theorem, without pretending that a literal universal source table has been materialized here. We count all gate factors and radii, give exact clocks, and construct degree-at-most-four sum-of-squares certificates for fixed finite source horizons, with exact natural witness fibers and explicit costs. A small fully materialized accepting machine, its complete raw particle orbit, vector figure, and exact replay tests accompany the proof. A separate, byte-preserved four-mass decidability report supplies the matching lower bound. These are mathematical proofs with independent computational and mathematical audits, not proof-assistant formalizations or claims of novelty or priority.

[write] Delivered as `Reversible_Binary_Five_Particle_Package.zip` (batch 82, archive 15; arrival `db37d18c8`), dated 3 October 2026; files placed by `bca6383e9` with the prefix given in Appendix A. This Report is the base of the merge: its text is printed in full, with its three `\input` files inlined.

4 Scope and theorem

We use $\mathbb{N} = \{0, 1, 2, \dots\}$. A binary configuration is an element $x \in \{0, 1\}^{\mathbb{Z}}$. Its support is $\text{supp}(x) = \{i : x_i = 1\}$; for finite support, its numerical mass is $M(x) = |\text{supp}(x)|$. A radius- R cellular automaton (CA) is given by one map $f : \{0, 1\}^{2R+1} \rightarrow \{0, 1\}$ used at every integer site and every time. Reversible means bijective on the *entire* full shift, not merely injective on encoded or finite configurations. We call it number-conserving when $M(Fx) = M(x)$ for every finite-support x . The all-zero configuration is fixed by the construction.

Theorem 4.1 (Effective reversible five-particle compiler). *Let C be a finite deterministic partial-injective two-natural-counter machine of the form in Section 5, with m controls, p nonzero-update branches, a zero-update branches and class cutoff J . Suppose its designated halt control has no outgoing branch. An effective procedure constructs:*

1. a binary, translation-equivariant, time-independent, finite-radius CA F that is bijective on $\{0, 1\}^{\mathbb{Z}}$ and preserves finite numerical mass;

2. a computable injection from source IDs into configurations with exactly five ones and zero elsewhere; and
3. one finite anchored binary word whose occurrence along the encoded forward orbit is equivalent to eventual reachability of the source halt control.

The factor count, radius bound, encoding diameter and exact source-edge clock are given explicitly below. No spatial checkerboard, external step counter, changing background or additional stored particle is used.

Corollary 4.2 (Universality by a fixed source). *There exists a fixed globally reversible, number-conserving binary CA of finite radius for which occurrence of a fixed finite anchored halt word is undecidable on a computable family of initial configurations containing exactly five ones.*

Proof. Choose a fixed universal deterministic Turing machine, apply the effective reversible two-counter simulation theorem of Morita [2], make the terminal-control normalization in Section 5, then apply Theorem 4.1. The resulting source, compiler parameters, CA and observer are fixed; only the effectively encoded finite input varies. $\square$

This is universality in an effective finite-input halting-observation sense. It is not a claim of intrinsic universality, real-time simulation, efficient input conversion, radius optimality, or a printed universal binary truth table. The explicit executable example in this release is finite and accepting, but is *not* universal. The universal reversible source is supplied by an existence-by-effective-construction theorem, rather than by an instantiated list of universal transition rows in this report.

[write] The literal universal reversible source that this paragraph says is absent is supplied in Part II: Report 16 instantiates it (Theorem 33.1) and feeds it to the compiler of Theorem 4.1, so the fixed automaton of Corollary 4.2 now has explicit parameters (Section 38).

5 Source machine and primary attribution

A source ID is $(q, c_0, c_1) \in Q \times \mathbb{N}^2$. Each branch e specifies source q_e , target q'_e , a selected counter $j_e \in \{0, 1\}$, an update $\Delta_e \in \{-1, 0, 1\}$ and a finite Boolean predicate $G_e(c_0, c_1)$ built from atoms $c_j = k$ and $c_j > k$, where $k \in \mathbb{N}$. The other counter is unchanged. The following are hypotheses on *all* source IDs:

1. enabled branches have natural-number outputs;
2. branch domains with a common source control are pairwise disjoint;
3. actual branch image domains with a common target control are pairwise disjoint.

The exact image predicate is

$$I_e(c') = [c' - \Delta_e \mathbf{e}_{j_e} \in \mathbb{N}^2] \wedge G_e(c' - \Delta_e \mathbf{e}_{j_e}). \quad (1)$$

It is computed from the branch; it is not an independently trusted guard. These hypotheses make the source transition relation a partial injection. An ID outside every branch domain is stuck. It counts as a halt observation only if its control is the designated h .

Choose a finite J such that every domain and image predicate is determined by each counter's class

$$\mathcal{K}_J = \{0, 1, \dots, J, \star\}, \quad \star = \{J + 1, J + 2, \dots\}.$$

Finite class enumeration verifies disjointness for predicates satisfying this cutoff promise. The parser uses a sufficient syntactic criterion that includes all atom thresholds and their updated-counter shifts. This criterion can reject logically redundant formulas even if a smaller cutoff would suffice; that is a deliberate validation restriction. Counter values are exact nonnegative integers. Floats and Boolean values masquerading as integers do not provide source IDs.

5.1 What is imported from Morita, and what is added here

Morita’s 1996 Definitions 2.1–2.3 (printed pp. 304–306) specify natural-counter IDs and reversible counter-machine transitions. Theorem 4.1 (p. 313) reduces reversible k -counter computation to two counters by an effective prime encoding. Theorem 4.2 (p. 319) simulates each deterministic Turing machine by a reversible two-counter machine [2]. Thus no uncounted history or work register remains beyond the final two counters and finite control.

Its row types translate into this interface as follows:

Row type	Domain on selected counter c	Update
zero test	$c = 0$	0
positive test	$c > 0$	0
increment	true	+1
decrement	$c > 0$	−1
no change	true	0

Here $J = 1$ suffices for all domains and exact images. For example, increment’s image is positive, while an increment restricted to zero has image exactly one. A reverse gate must use these images, not blindly reuse a forward test.

The cited definition designates a final control but does not explicitly ban outgoing rows from it. *Our normalization* deletes all such rows. This restricts a partial injection and preserves first reachability of that control. It does not claim to preserve post-halt behavior. If a predecessor-free initial control is wanted, Morita–Imai 2001, Lemma 3.5 (printed pp. 245–246), supplies a reversible two-counter normal form with no row entering its initial control [3]. Apply that first, then delete final-control exits. This extra condition is used for exact finite reflection cycles and the clean target construction, but not for the forward halt-pattern equivalence.

Prime encoding can be expensive in time and value. The input conversion need only be computable. No numerical values of m, p, a for a literal universal reversible table are claimed here, and no nonreversible source table from a previous report is silently substituted.

[write] Report 16 (Part II) gives such numerical values for one literal universal reversible table in this schema: $m = 122,622$ controls, 141,561 branches and $J = 0$. The nonreversible 8,408-row source of the second half of this Part (Report 14) is not that table, and Report 16 does not use it.

6 An isolated swap on every binary input

The all-input lemma is the reason malformed configurations cannot undermine global reversibility. Fix $B \geq 1$ and two finite supports $P, Q \subseteq [-B, B]$ with equal nonzero cardinality k . Suppose P and Q are not translates. Their coordinate zero is an assigned key, not necessarily the leftmost occupied site. A position u is a *raw key* in x when the ones in $[u - B, u + B]$ are exactly $u + P$ or exactly $u + Q$. Let a local, translation-covariant contextual predicate read within radius T of the key, and assume it has the *same truth value at paired endpoints* when outside bits are fixed. Set

$$L_B = 3B + 1, \quad M_B = \max(L_B + B, T + B) + 1, \quad r_B = M_B + 2B. \quad (2)$$

A raw key u is eligible if:

1. its only ones in $[u - L_B, u + L_B]$ are its displayed k particles;
2. no other raw key lies at distance at most M_B ; and
3. the contextual predicate holds.

Swap $u + P$ and $u + Q$ at all eligible keys simultaneously and leave every other bit unchanged. Denote this map by $\text{Swap}(P, Q; B, G, T)$.

Lemma 6.1 (All-input isolated-swap lemma). *The displayed map is an involutive binary CA of radius at most r_B . It preserves numerical mass on every finite configuration. These statements hold without any density, finite-support or well-formedness promise on the input for the involution and locality claims.*

Proof. Distinct eligible keys are farther than M_B apart. Their radius- B cores are disjoint. This makes even an infinite family of simultaneous rewrites well-defined. No other active core intersects the radius- L_B isolation halo or radius- T contextual window of an active key, since intersection would force key distance at most $L_B + B$ or $T + B$.

Write y for the rewritten configuration. We first prove that the *entire raw-key set*, not merely the active subset, is unchanged. A raw-test window that changes intersects the rewritten core of some active key u . Its key v has $|v - u| \leq 2B$, so its entire window is contained in $[u - 3B, u + 3B]$. In x this larger interval has just $u + P$ or $u + Q$ as its k ones. In y it has precisely the other support. If the window at v is raw at either endpoint, its k ones must therefore be the *entire* support of the active pattern at that endpoint. Nontranslation of P, Q forces its label to be that of the active pattern. A nonempty finite set cannot equal a nonzero translate of itself, so $v = u$. Every active key remains raw after its own rewrite. All other raw-test windows that could change are nonraw on both sides. This proves raw-key invariance, including changes involving a required zero in a matching window.

Consequently, the no-other-raw-key condition is invariant. An active key retains isolation because its only changed particles are its own equal-mass core. It retains the contextual condition by the paired-endpoint assumption and separation from all other active cores. An inactive raw key that is blocked by another raw key stays blocked. An inactive raw key with no such blocker is farther than M_B from every active key, so none of its code, isolation or context tests can change. Thus eligibility is unchanged at every key. A second application rewrites exactly the same disjoint cores and restores their original patterns: $\text{Swap}^2(x) = x$ for every x .

For an output site i , only keys within B can write at i . A key's eligibility is determined by its radius- $(M_B + B)$ input window, which covers the raw-key exclusion, isolation and context tests. Radius $M_B + 2B$ about i therefore suffices. Every test is in relative coordinates, proving translation equivariance. Finally, for finite input there are finitely many disjoint active cores, and each replaces k ones by k ones. This proves mass conservation. $\square$

The guard requirement is genuinely cross-endpoint. Giving the P endpoint a constant true guard and the Q endpoint a constant false guard would break the involution, although each constant is individually insensitive to bits. Every contextual gate below uses one identical predicate of unchanged exterior bits at both endpoints, so it meets the stronger requirement directly.

A finite ordered composition of these gates is a full-shift reversible, number-conserving CA. Its inverse applies the same involutions in reverse order. The ordered product itself need not be an involution.

7 Five-particle geometry

There are $m + 2p$ unsigned head modes: H_q for each source control, and O_e, l_e for each nonzero branch. Give each signed mode u^+, u^- a *distinct* gap $d(u^\pm)$ in $\{1, \dots, D\}$, where

$$\begin{aligned} D &= 2m + 4p, & S &= 2D + 2, & B_2 &= D + 1, \\ L &= 3D + 4, & B_3 &= 4D + 5, & Z &= 10B_3 + 10 + 2J = 40D + 60 + 2J. \end{aligned} \quad (3)$$

A signed head u^σ at anchor x is the pair $\{x, x + d(u^\sigma)\}$. The anchor is always its leftmost particle, even during leftward travel.

The home encoding, with origin zero, is

$$\text{enc}(q, c_0, c_1) = \{-(Z + c_0), 0, Z + c_1, S, S + d(H_q^+)\}. \quad (4)$$

It has exactly five distinct occupied sites. Three singleton markers encode the left endpoint, origin, and right endpoint. The endpoints carry counter values by their distances to the origin. The remaining close pair encodes all control and sign information. The singleton particles have no hidden colors. The initial diameter is $2Z + c_0 + c_1$.

On the admissible microstates below, marker separations are at least Z . All head-to-marker distances are at least $S - D = D + 2$, while the head's internal gap is at most D . Thus the head is the unique close pair. Removing it leaves three markers whose middle member is the origin. This proves effective geometric decoding and disambiguates roles with only binary sites. In particular, signed mode information is paid for in pair spacing and radius, not in additional particles or a larger alphabet.

8 The partial micrograph and its exact clock

For a nonzero branch e , write $v = -1$ for the left counter and $v = +1$ for the right. Let o be the origin, m the selected endpoint, and c its old counter value. Replace its source edge by the following finite path of unsigned microstates:

1. **Dispatch.** Replace a home H_{q_e} head at $o + S$ by an O_e head at $o + vS$, under G_e .
2. **Outbound travel.** Move its anchor by v at each microedge, from $o + vS$ to $m - vS$.
3. **Endpoint interaction.** Move the selected endpoint to $m' = m + v\Delta_e$ and change the head to l_e at $m' - vS$.
4. **Return travel.** Move its anchor by $-v$ at each microedge until $o + vS$.
5. **Commit.** Change to $H_{q'_e}$ at $o + S$, under the exact image predicate I_e evaluated on the post-counters.

A zero-update branch is one directly guarded home-to-home microedge.

Definition 8.1 (Admissible microstates). *The unsigned admissible set consists of every home ID in $Q \times \mathbb{N}^2$ and only the strict intermediate configurations on a displayed subdivision of an enabled source edge e, c satisfying $G_e(c)$. The doubled admissible set contains a plus and minus copy of exactly this set.*

Arbitrary geometrically plausible outbound or return packets are *not* included merely because their spacings look reasonable. Branch-domain and exact-image compatibility are part of admissibility. Global CA reversibility will not depend on this restriction, but simulation closure does.

Lemma 8.2 (Injective subdivision and time). *The micrograph on admissible unsigned states is a partial injection. A nonzero branch at selected counter c has exactly*

$$\tau_e(c) = 3 + 2(Z + c) + \Delta_e - 4S \quad (5)$$

microedges; a zero-update branch has one.

Proof. Distinct branch labels distinguish intermediate paths. An outbound mode retains old counters and branch label; a return mode retains the post-counters and branch label, from which the old counters are recovered. Its anchor fixes its unique position within that corridor. Home successors are unique by source-domain disjointness. Home predecessors are unique by disjointness of actual branch images. Therefore subdivisions neither branch nor merge improperly. Their outbound and return travel lengths are respectively $Z + c - 2S$ and $Z + c + \Delta_e - 2S$. Both are nonnegative with a large strict margin by (3). Adding dispatch, endpoint interaction and commit gives (5). $\square$

For motion of direction $w \in \{-1, +1\}$ near a marker m , a marker behind the head has anchor $x = m + wt$, with travel allowed for $t \geq S$. A marker ahead has $x = m - wt$, with travel allowed for $t > S$; at $t = S$ the next edge is the appropriate interaction. There is no travel edge from behind distance $S - 1$ into a dispatched or reversed state. These endpoint conventions are essential to the matching argument.

9 The edge matching as concrete pattern swaps

For each unsigned microedge $z \rightarrow z'$, we want the transposition $z^+ \leftrightarrow z'^-$. The following finite list of individual isolated swaps is composed once, in any one fixed order, to form the edge block E . The fixed order is part of the compiled local rule.

9.1 Free flight

For each moving mode u with direction w , use a two-particle gate with $B = B_2$, $T = 0$ and supports

$$P = \{0, d(u^+)\}, \quad Q = \{w, w + d(u^-)\}. \quad (6)$$

Its invariant key is the *predecessor* anchor x . It is eligible on an admissible predecessor exactly when no marker lies within $[x - L, x + L]$. The gate's raw key is unique because any raw pair must be the unique close pair. The reverse endpoint still checks isolation at x , not at its new anchor $x + w$.

9.2 Contact travel

For each moving u of direction w , use three-particle gates keyed at the contacted marker zero, with $B = B_3$, $T = 0$:

$$\{0, wt, wt + d(u^+)\} \longleftrightarrow \{0, w(t+1), w(t+1) + d(u^-)\}, \quad S \leq t \leq L; \quad (7)$$

$$\{0, -wt, -wt + d(u^+)\} \longleftrightarrow \{0, -w(t-1), -w(t-1) + d(u^-)\}, \quad S+1 \leq t \leq L. \quad (8)$$

These cover the complement of free flight on valid predecessors near a marker. All integer t in the displayed ranges give separate finite gate factors.

The free/contact boundary, both directions. Behind a marker, departure from predecessor distance L uses (7), ending at distance $L+1$. The potential inverse of (6) still tests the predecessor key at distance L , and therefore sees the marker and is blocked. Ahead of a marker, arrival from predecessor distance $L+1$ uses (6), ending at distance L . There is no gate (8) with predecessor distance $L+1$, so it cannot compete. At behind distance S , an outgoing travel edge exists but no incoming travel edge from $S-1$ does; dispatch or endpoint reversal supplies that incoming edge. At ahead distance S , travel has no outgoing edge; endpoint interaction or commit supplies it. These statements remain true for $w = -1$ because (6) preserves the predecessor key rather than redefining it by the sign of travel.

9.3 Dispatch, endpoint and commit

Each nonzero branch supplies three triple gates, with $B = B_3$:

$$\{0, S, S + d(H_{q_e}^+)\} \longleftrightarrow \{0, vS, vS + d(O_e^-)\}, \quad (9)$$

$$\{0, -vS, -vS + d(O_e^+)\} \longleftrightarrow \{v\Delta_e, v\Delta_e - vS, v\Delta_e - vS + d(I_e^-)\}, \quad (10)$$

$$\{0, vS, vS + d(I_e^+)\} \longleftrightarrow \{0, S, S + d(H_{q_e}^-)\}. \quad (11)$$

The dispatch gate uses G_e , the commit gate uses I_e , and the endpoint gate has no context. The endpoint gate's invariant key is the *old* endpoint coordinate, including when its reverse endpoint is inspected.

For a zero-update branch use the triple gate

$$\{0, S, S + d(H_{q_e}^+)\} \longleftrightarrow \{0, S, S + d(H_{q_e}^-)\}, \quad (12)$$

with the same source/image predicate $G_e = I_e$.

9.4 Reading bounded counter classes without storing anything

At a home gate's origin key, read the bits at $-(Z+k)$ and $Z+k$ for $0 \leq k \leq J$. In an admissible state one present bit determines its finite class; all-zero reads indicate $\star$. If more than one bit appears on a side on a malformed input, define the guard false. This is a finite Boolean word test of radius $T = Z + J$.

All read positions are outside the swapped support. At each of (9), (11) and (12), the same origin key, same exterior bits, and same predicate are used at both endpoints. Hence the contextual gate satisfies Lemma 6.1 on the full shift, not just the intended data.

Every displayed triple has three distinct sites, one uniquely close pair of gap at most D , and other gaps greater than D . Plus and minus gaps differ, so the two supports are not translates. Their extremes lie in $[-B_3, B_3]$; in particular, the largest contact departure endpoint is $L+1+D = B_3$. Other markers are outside its particle-isolation halo because $Z > 3B_3 + 1$, with margin even for the old-key displacement in (10). Any raw triple must contain the unique close pair and its nearby marker. Its gap identifies the endpoint label, and then its anchor identifies the key. Thus there is only one raw key on an admissible state. A large contextual raw-key-exclusion radius cannot accidentally veto the intended gate.

Lemma 9.1 (Matching and closure). *On the doubled admissible set, E acts exactly as the matching that exchanges z^+ with z'^- for each unsigned microedge $z \rightarrow z'$, and fixes unmatched vertices. No order choice among its fixed gate factors changes this action.*

Proof. A plus state belongs only to its outgoing microedge; a minus state belongs only to its incoming microedge. Equations (6)–(12) realize precisely those edges, with the travel-boundary

cases above eliminating overlap and gaps. At homes, plus incidence is unique by domain disjointness and minus incidence by exact image disjointness. Conversely, a minus home passing I_e has a unique natural predecessor $c = c' - \Delta_e \mathbf{e}_{j_e}$ satisfying $G_e(c)$ by (1). Its inverse commit or direct gate therefore enters an *included* subdivision, not spurious moving geometry. Forward dispatch similarly enters an enabled subdivision. Branch labels preserve intermediate uniqueness. Thus every gate stays in the admissible set and realizes its one matching edge there. Once one factor moves a vertex to its partner, no other matching edge is incident at the partner, so subsequent factors do nothing to it. $\square$

The matching claim is restricted to admissible configurations. It is not a claim that E itself is an involution on every malformed configuration. Full-shift reversibility follows instead from the independently involutive factors, which require no matching hypothesis.

10 Phase flip, reflection and observation

The signed phase block P uses the following factors. For each moving u , include the pair gate

$$\{0, d(u^+)\} \longleftrightarrow \{0, d(u^-)\} \quad (B_2, T = 0). \quad (13)$$

For each moving u , each $s \in \{-1, +1\}$ and $S \leq t \leq L$, include

$$\{0, st, st + d(u^+)\} \longleftrightarrow \{0, st, st + d(u^-)\} \quad (B_3, T = 0). \quad (14)$$

For each $q \in Q$, include the home triple

$$\{0, S, S + d(\mathbf{H}_q^+)\} \longleftrightarrow \{0, S, S + d(\mathbf{H}_q^-)\} \quad (B_3, T = 0). \quad (15)$$

On an admissible state, precisely one phase gate is active. At a moving anchor with no marker within L , it is the pair gate. Otherwise there is one nearby marker at a distance in $[S, L]$, and exactly one contact phase gate applies. At a home the origin supplies its home phase gate. A phase flip keeps the anchor fixed, so it cannot change this classification. The fixed-order block P therefore acts as sign reversal on admissible states.

Define *one CA time step* to be

$$F = P \circ E. \quad (16)$$

This is one fixed composition of finitely many translation-equivariant local rules. Composition of radius- r_i rules has radius at most $\sum_i r_i$. Evaluating the factors internally is an algorithm for the single local map f ; it is not an externally alternating CA, a spatial phase convention or an extra clock layer. Its full-shift inverse is the reverse list of the same involutions. Thus Lemma 6.1 proves full-shift bijectivity and finite mass conservation independently of simulation validity.

On an admissible plus vertex with successor, E moves to its successor's minus copy and P restores plus. On a minus vertex with predecessor, E moves to its predecessor's plus copy and P restores minus. At a missing outgoing or incoming edge, E fixes the vertex and P flips its sign in one step. This proves the desired reversible reflection dynamics.

Proposition 10.1 (Exact reflection cycle). *If a starting source ID has no predecessor and its forward computation reaches terminal halt after T microedges, the encoded orbit has exactly $2T + 2$ states, visiting the full forward history, then the full reverse history, then returning to its initial plus state.*

Proof. A finite partial-injective path from a predecessor-free start to a terminal vertex has $T + 1$ distinct unsigned vertices. Plus and minus copies are distinct because their gaps differ. The forward and backward traversals take T steps each; the two boundary sign flips each take one. This visits each of the $2T + 2$ signed vertices once before returning. $\square$

In particular, a reversible orbit does not enter a separate periodic attractor or freeze at halt. Its backward history is part of its own cycle.

The observer is the anchored binary word on $[0, S + D]$ with ones exactly at

$$\{0, S, S + d(H_h^+)\}, \quad (17)$$

and zeros at all other sites in that interval. Its length is $3D + 3$. Both endpoints lie outside the word. The unique head gap identifies the halt plus home. No moving head or other home mode can give that gap. The interval's origin and fixed head position force the home geometry on admissible states.

If the forward source run reaches h , the observer occurs. For the converse, a different stuck home may reflect, but a backward traversal cannot newly reach an h home: that would require an outgoing microedge from h , and h is terminal. Reflection only retraces the same path component. Therefore the observer occurs along the full encoded orbit if and only if the forward source computation reaches h . The negative halt code is different. This also explains why terminal-control normalization is required for the stated unrestricted-in-time observation equivalence.

11 The exact resource ledger

All factors use alphabet $\{0, 1\}$; all encoded microstates have mass five. Unoccupied guard-reading positions store no data. With constants (3), the factor counts and individual radius bounds are:

Factor category	Number	Radius bound
Pair gates	$4p$	$r_2 = 6D + 8$
Unguarded triple gates	$8pD + 23p + m$	$r_3 = 24D + 32$
Contextual triple gates	$2p + a$	$r_c = Z + J + 12D + 16$
All factors	$8pD + 29p + m + a$	sum below

To check the count, each of $2p$ moving modes has one free-flight and one pair-phase gate. Its contact travel contributes $(L - S + 1) + (L - S) = 2D + 5$ triples; contact phase contributes $2(L - S + 1) = 2D + 6$. Add p unguarded endpoint triples, m home-phase triples, $2p$ contextual dispatch/commit triples, and a contextual direct triples. This gives every row above.

For $T = 0$, (2) gives radius $6B + 2$, which yields r_2, r_3 after substitution. For contextual triples, $T = Z + J$ dominates and the radius is $T + 3B_3 + 1 = r_c$. Therefore a valid explicit composition radius is

$$\begin{aligned} R = & 4p(6D + 8) + (8pD + 23p + m)(24D + 32) \\ & + (2p + a)(Z + J + 12D + 16). \end{aligned} \quad (18)$$

This is a counted upper bound, not a minimum radius. A literal local truth table has 2^{2R+1} binary entries. It can be generated by finite composition of the factor algorithms, but writing it explicitly would typically be far larger than the gate description. Compilation is effective without any claim that this truth-table expansion is practical.

For a source path of H branches with old selected values c_{t,j_t} , its first halt time is

$$T = \sum_{t=0}^{H-1} \begin{cases} 1, & \Delta_{e_t} = 0, \\ 3 + 2Z + \Delta_{e_t} - 4S + 2c_{t,j_t}, & \Delta_{e_t} \neq 0. \end{cases} \quad (19)$$

A certificate's source horizon is H , not this counter-dependent CA time. After at most H source steps, both counters are at most their initial value plus H , their sum is at most the initial sum plus H , and the entire subdivided path remains inside a span of at most

$$2Z + c_{0,0} + c_{0,1} + H. \quad (20)$$

The head stays between endpoints throughout. These are value/diameter bounds; they are not bit-complexity estimates.

12 Counted finite-horizon certificates

This algebraic construction needs deterministic source semantics, natural outputs and a terminal halt; its proof does not need reversibility. The additional image-disjointness validation connects its source interface to the reversible CA compiler. Fix natural initial counters $c_{0,0}, c_{0,1}$, a control q_0 , and a source horizon H . The horizon is a *compiler argument*: a separate finite polynomial is generated for each H . It is not an unbounded variable inside one fixed polynomial.

12.1 Disjoint finite class normalization

For each original branch, enumerate all $(J+2)^2$ pairs of classes from $\mathcal{K}_J$. Retain (e, α_0, α_1) exactly when its guard is true on that class cell. Call these retained normalized branches β . Write B for their count, r for the total number of tail-coordinate occurrences among them, and $P_{\neq 0}$ for their nonzero-update count. Then

$$B \leq (p+a)(J+2)^2, \quad 0 \leq r \leq 2B, \quad P_{\neq 0} \leq B. \quad (21)$$

Also $B \leq (m-1)(J+2)^2$ by determinism and terminal halt. Exactly $(p+a)(J+2)^2$ original-branch/cell pairs are examined. Boolean formula length is an additional finite evaluation cost. Unsatisfiable rows retain no cells.

Each normalized branch retains the original source and target controls, selected counter, update and duration. This is a partition of a branch's domain, not a new sequence of instructions. It neither increases the source horizon nor changes m, p, a, D, S, Z . Full cells avoid duplicate witnesses from formulas such as $A \vee A$, which would be easy to introduce by a careless disjunctive expansion.

Use distinct integer control codes $0, \dots, m-1$, with source and target codes s_β, t_β . Let $\delta_{\beta,k}$ denote its counter- k update, zero on the unselected counter. For $H \geq 1$, introduce natural variables

- $e_{t,\beta}$ for $0 \leq t < H$ and all B normalized branches;
- postcounters $c_{t+1,0}, c_{t+1,1}$ at each layer;
- a slack $z_{t,\beta,k}$ exactly for each tail coordinate $\alpha_{\beta,k} = \star$.

Initial counters and control are external parameters or specialized constants, and are not counted as witness variables.

12.2 Every residual explicitly

The polynomial is the sum of squares of the following integer polynomials:

$$A_t = \sum_{\beta} e_{t,\beta} - 1, \quad (22)$$

$$Q_0 = \sum_{\beta} s_{\beta} e_{0,\beta} - q_0, \quad (23)$$

$$Q_t = \sum_{\beta} s_{\beta} e_{t,\beta} - \sum_{\beta} t_{\beta} e_{t-1,\beta} \quad (t > 0), \quad (24)$$

$$U_{t,k} = c_{t+1,k} - c_{t,k} - \sum_{\beta} \delta_{\beta,k} e_{t,\beta}, \quad (25)$$

$$E_{t,\beta,k} = e_{t,\beta} (c_{t,k} - \alpha_{\beta,k}) \quad (\alpha_{\beta,k} \neq \star), \quad (26)$$

$$T_{t,\beta,k} = e_{t,\beta} (c_{t,k} - J - 1) - z_{t,\beta,k} \quad (\alpha_{\beta,k} = \star), \quad (27)$$

$$F_H = \sum_{\beta} t_{\beta} e_{H-1,\beta} - h. \quad (28)$$

There is exactly one class residual, either (26) or (27), for every (t, β, k) . Let $\mathcal{P}_H$ be the sum of the squares of all these residuals. Each residual has degree at most two, so $\deg \mathcal{P}_H \leq 4$. This is an ordinary polynomial with integer coefficients; sum-of-squares notation is a finite representation of it, not an extra unbounded predicate. There are no exponentiation atoms, floor operations or hidden integrality tests.

Theorem 12.1 (Exact complete natural witness fiber). *For fixed natural initial input, the complete zero fiber of $\mathcal{P}_H$ in its natural witness domain is empty or a singleton. It is a singleton exactly when the source first reaches its designated halt control after exactly H original source instructions.*

Proof. The sum of squares is zero exactly when each residual vanishes. Over $\mathbb{N}$, $A_t = 0$ forces exactly one selector to be one and all others zero. The state rows select a branch at the current control, and the counter rows impose its exact update. If its selected cell is exact, (26) enforces that old counter value. If its selected cell is a tail, (27) forces $z = c_{t,k} - J - 1 \geq 0$, giving precisely the required class. Inactive tail rows force $z = 0$, so even their coordinates are uniquely determined. Therefore all selected rows are genuinely enabled, and all updates remain natural. The last residual enforces terminal halt. Since halt has no outgoing branch, it cannot have been reached at an earlier layer of an H -step legal trace.

Conversely, the deterministic first-halt trace specifies its one-hot selectors, postcounters, selected tail slacks, and zero inactive slacks, and these satisfy every row. Induction using domain disjointness makes the normalized branch and every coordinate unique at every layer. This proves uniqueness of the *entire* tuple, not just uniqueness after discarding auxiliary variables. $\square$

For $H = 0$, use $(q_0 - h)^2$ with no witness variables. The empty tuple is its unique zero exactly when $q_0 = h$. If $H > 0$ but $B = 0$, each one-hot residual is -1 , so no witness exists, as required. Other stuck controls never count as halt.

12.3 Paying for the exact CA clock

For a nonzero original branch define

$$\kappa_e = 3 + 2Z + \Delta_e - 4S = 72D + 115 + 4J + \Delta_e.$$

Its duration is $\kappa_e + 2c_{t,j_e}$; a zero-update row's duration is one. Using *the original compiler constants*, add a natural variable Θ and the residual

$$K = \Theta - \sum_{t,\beta} e_{t,\beta} \tau_{\text{original}(\beta)}(c_{t,j_\beta}). \quad (29)$$

This adds one witness and one squared residual, keeps degree at most four, and uniquely forces Θ . The CA simulation theorem identifies it with the first anchored plus-halt observation time. The phase reflection *after* halt is not charged. Counting individual factors within a CA step would be a different and incorrect clock convention here.

Alternatively, specialize Θ as a natural external input; then only the square is added. At $H = 0$, the unsimplified clock row is simply Θ , so a clocked witness version has one variable and two residual slots. A deadline N can be imposed with one more natural slack d and one square $(N - \Theta - d)^2$, uniquely forcing $d = N - \Theta$. This still represents fixed source horizon H . Although positive step durations imply $H \leq \Theta \leq N$ on accepting runs, combining all horizons $H \leq N$ requires an additional explicitly charged construction.

12.4 Exact counts, expansion cost and height

For $H \geq 1$, the core witness and residual counts are

$$\boxed{V = H(B + 2 + r)}, \quad \boxed{S_{\text{SOS}} = H(4 + 2B) + 1}, \quad \deg \mathcal{P}_H \leq 4. \quad (30)$$

Residual slots are counted even when a specialized input makes a row identically zero. The clock adds one variable and one slot. Eliminating state variables by selector-coded controls is already reflected in these counts.

Before collecting terms or specializing constants, a bound on the written residual monomial slots is

$$H(7B + P_{\neq 0} + r + 5) + 2. \quad (31)$$

Indeed the one-hot rows use $H(B + 1)$ slots; state rows use $B + 1$ initially and $2B$ thereafter; the two update rows together use at most $P_{\neq 0} + 4$ per layer; class rows use at most $4B + r$ per layer; and the terminal row uses $B + 1$. The clock adds at most $1 + H(B + P_{\neq 0})$ slots. First-layer external input coordinates are counted as atomic terms; any polynomial substitution for them must pay its own expansion cost.

Let L_i be the actual collected term count in residual i . Fully expanding the sum of squares produces at most $\sum_i L_i^2$ ordered term occurrences before final collection. Both $\sum_i L_i$ and $\sum_i L_i^2$ are emitted, and the emitter can also compute an exactly collected expanded polynomial. Thus squaring is explicitly costed, not treated as free.

Write $M = \max(c_{0,0}, c_{0,1})$ and $\kappa_{\max} = \max(1, \max_{\Delta_e \neq 0} \kappa_e)$. At the unique valid witness,

$$\begin{aligned} 0 &\leq e_{t,\beta} \leq 1, & 0 &\leq c_{t+1,k} \leq M + H, \\ 0 &\leq z_{t,\beta,k} \leq M + H, & 0 &\leq \Theta \leq H(\kappa_{\max} + 2M) + H(H - 1). \end{aligned} \quad (32)$$

The clock bound sums $c_{t,j} \leq M + t$. All non-clock coordinates therefore need at most $\lceil \log_2(\max(1, M + H) + 1) \rceil$ bits each in ordinary unsigned binary. With the clock, replace that height by the maximum of $1, M + H, H(\kappa_{\max} + 2M) + H(H - 1)$. These are witness-value and representation bounds, not the complexity of finding a witness or encoding a universal input. The conditional CA diameter bound is (20).

With fixed initial constants, a safe bound on collected core residual coefficients is

$$C_0 = \max(2, m - 1, J + 1, M + J + 1).$$

With the clock use $C_1 = \max(C_0, \kappa_{\max} + 2M, 2)$, and include N if a deadline is present. A deliberately coarse expanded coefficient bound is $(\sum_i L_i^2)C_1^2$, by the triangle inequality over all ordered products. This includes possible coefficient collection and cancellation. For symbolic initial counters, their monomials must be counted explicitly and M is removed from the coefficient bound, while value-height bounds still depend on the evaluated natural input.

12.5 The paid nonnegative-real variant

Natural selectors summing to one are automatically one-hot; real selectors are not. To obtain the same complete fiber over the nonnegative real orthant, add one residual per layer:

$$N_t = \sum_{\beta} e_{t,\beta}^2 - 1. \quad (33)$$

When $A_t = N_t = 0$ and all selectors are nonnegative, $\sum_{\beta < \gamma} e_{t,\beta} e_{t,\gamma} = 0$. Every summand is nonnegative, so exactly one selector is one. From natural initial counters, the update rows inductively force natural postcounters, then all tail slacks and the clock are natural too. Hence the entire orthant zero fiber is exactly the same empty-or-singleton natural tuple.

The price is H additional squared residual slots and $H(B + 1)$ raw monomial slots, with no additional variables and degree still at most four. This is neither an unrestricted-real theorem nor a theorem for arbitrary real initial counter values. Exact software checks use integers and rational numbers; the mathematical orthant argument applies to all nonnegative reals.

[write] The same norm row is used by Report 14 for its own certificates (equation (29.1) and Theorem 29.1); it is printed once here and once there because the two certificates have different selectors. For degree-two certificates the collection uses a different device, products forcing a single selector (quadratic-orthant-certificates, Lemma `qoc:rn:lem:gates`).

12.6 Materialized certificate of the small source

For the one-increment source in Section 14, $J = 1$ gives nine retained cells, six tail occurrences and nine normalized nonzero branches: $B = 9, r = 6, P_{\neq 0} = 9$. At $H = 1$ the core has 17 variables and 23 squared residual slots. Adding the clock gives 18 variables and 24 slots, with its unique witness forcing $\Theta = 696$. The paid orthant version has the same 18 variables and 25 slots. The JSON artifacts list every variable, coefficient and monomial and the full exact witness, so this is an emitted finite polynomial rather than a notation-only example.

No claim of an unbounded single-fold Diophantine representation follows by varying H outside the polynomial. Likewise, ordinary zero/positive-only counter certificates cannot be copied unchanged for arbitrary Boolean guards: the explicit disjoint cell normalization and all rH uniquely forced slack coordinates are essential to the present theorem and counts.

13 Clean exact targets without another particle

Suppose the source's initial *control* q_0 has no incoming branch at any counter values, and h has no outgoing branch. A predecessor-free *particular input ID* is insufficient for the unconditional bridge used below. The control-wide normalization discussed in Section 5 is available for the classical universal source.

Make controls $F(q), B(q)$ for every $q \in Q$ and a fresh terminal control H_* . For every original branch include

$$\begin{array}{ll} F(q_e) \longrightarrow F(q'_e) & \text{with update } \Delta_e \text{ and guard } G_e, \\ B(q'_e) \longrightarrow B(q_e) & \text{with update } -\Delta_e \text{ and guard } I_e. \end{array}$$

Add two zero-update true-guard bridges

$$F(h) \longrightarrow B(h), \quad B(q_0) \longrightarrow H_*. \quad (34)$$

The reverse domains are exactly the original images, and reverse images are exactly the original domains. Their disjointness conditions are therefore exchanged. The first bridge fills a missing outgoing slot at $F(h)$ and missing incoming slot at $B(h)$, using terminality of h . The second fills the missing outgoing slot at $B(q_0)$ and enters an otherwise unused control, using the control-wide no-incoming hypothesis. Thus the wrapper is deterministic and partial-injective on *all* natural IDs. Its initial control has no incoming row and H_* has no exit. Its counts are

$$m' = 2m + 1, \quad p' = 2p, \quad a' = 2a + 2. \quad (35)$$

The same semantic class cutoff J suffices because reverse domains/images are the old images/domains. A literal guard syntax may need simplification or a conservatively enlarged parser cutoff, which must then be reflected in the compiler constants; no unchecked image predicate is accepted as data. Write J_* for the actual valid serialized cutoff, using $J_* = J$ when the serialization meets the original cutoff.

On a halting input, the source runs forward to h , takes its bridge, retraces all work backward and takes the second bridge into $(H_*, c_{0,0}, c_{0,1})$. The original initial control cannot recur internally because it has no incoming row. The counters are restored exactly, without augmentation. If the original forward computation does not reach h , this wrapped forward computation cannot reach H_* . Apply the five-particle compiler to this enlarged table, *recomputing* all constants from m', p', a', J_* .

Corollary 13.1 (Exact finite-target reachability). *For a fixed compiler instance built from a universal source in the stated normal form, exact finite-configuration reachability is undecidable on pairs of five-particle configurations. The designated target for input $(c_{0,0}, c_{0,1})$ is*

$$\{-(Z + c_{0,0}), 0, Z + c_{0,1}, S, S + d(H_{H_*}^+)\}. \quad (36)$$

It is a computable input-dependent target, not a single constant target shared by every input.

Proof. The wrapper reaches its clean terminal ID exactly when the original input halts. The CA reaches the displayed plus-home configuration on that forward path. Conversely, later sign reflection cannot create a new terminal home that the forward path never visited, by Section 10. $\square$

[write] **Credit.** The forward copy, inverse copy and two bridges above are the clean wrapper of signal-machine-collision-certificates source 15 (Report 10 of this series; its equations `smc:ct:eq:copies` and `smc:ct:eq:bridges`, Lemma `smc:ct:lem:syntax`, Theorem `smc:ct:thm:restore`, and its `§smc:ct:sec:fresh` on the fresh-entry condition), here generalized to guarded sources, whose reverse guard is the exact image predicate (1). The construction is Bennett's uncomputation [8]. Report 15 cites neither; the clock (37) corresponds to that report's Theorem `smc:ct:thm:clock`, and the zero-set bijection of the next subsection to its Theorem `smc:ct:thm:lift`. No priority is claimed for the wrapper.

Let Θ' be the physical microedge count of the original forward computation measured with the *enlarged* compiler constants. For a nonzero step,

$$3 + 2(Z + c) + \Delta - 4S = 3 + 2(Z + c + \Delta) - \Delta - 4S.$$

Forward and reverse durations are equal, as are zero-update durations. Consequently the first clean-target time and full CA reflection period are

$$T_* = 2\Theta' + 2, \quad \Pi = 4\Theta' + 6. \quad (37)$$

For H original source instructions the wrapped source has $2H + 2$ instructions. The bounded certificate may be applied directly to that wrapped table and horizon, paying for its own normalized counts; alternatively an explicit rescaled clock of the original forward trace must use the enlarged constants, not its old geometry.

13.1 Full-history and projected clean certificates

Applying Section 12 to the validated wrapped table at source horizon $2H + 2$ gives

$$V_* = (2H + 2)(B_* + 2 + r_*), \quad S_* = (2H + 2)(4 + 2B_*) + 1.$$

The exact full clock adds one variable and one square; the paid orthant version adds $2H + 2$ squares. Final-counter restoration follows from the wrapper semantics. If one wants the target equalities visible as polynomial rows, add $c_{2H+2,k} - c_{0,k}$ for $k = 0, 1$, paying two more squares and no variables. They are redundant on the wrapper's zero set, and are not counted in the bundled core.

There is also a theorem-level forward-only option. Keep the original $\mathcal{P}_H$ and its class normalization, and append a natural variable T_* with the single residual

$$T_* - 2 \sum_{t,\beta} e_{t,\beta} \tau_{\text{original}(\beta)}^* (c_{t,j_\beta}) - 2,$$

where τ^* uses the enlarged geometry. This costs one variable, one square and at most $2 + H(B + P_{\neq 0})$ raw monomial slots, with degree at most four. For $H = 0$ it is simply $T_* - 2$. An accepting forward zero uniquely reconstructs the reverse trace, reverse cells and reverse slacks. Conversely, a wrapped zero projects to its forward trace. This is a bijective correspondence *on zero sets*; it is not equality of polynomials away from their zeros, nor a globally polynomial reverse-witness map. For example, a forward tail decrement can land in exact class J , so reverse selectors must be recomputed from the actual restored trace. The executable emitter supplies the full-history option; the projected formula here is not presented as an implemented command-line feature.

The fully materialized clean example uses the one-increment source and its inverse decrement, with two bridges. It has $m' = 5, p' = 2, a' = 2, J = 1, D = 18, S = 38, Z = 782$, 353 factors, radius bound 164314 and observer length 57. At input $(0, 0)$ its supports are

$$\text{initial } \{-782, 0, 38, 39, 782\}, \quad \text{target } \{-782, 0, 38, 47, 782\}.$$

Here $\Theta' = 1416$, first target time is 2834, and the checked full reflection period is 5670. Its complete raw trace is included. Here the full wrapped normalization has $B_* = 33, r_* = 23, P_{\neq 0,*} = 15$. At its source horizon four, the core has 232 witnesses and 281 squared residuals; the emitted clocked version has 233 witnesses and 282 squares, and the paid orthant version has 286 squares. Both emitted exact witnesses force clock 2834.

A limit on periodicity conclusions. A predecessor-free start returns to itself under reflected dynamics exactly when its unsigned forward path is finite; an infinite partial-injective path from such a start cannot repeat a vertex. But a finite path may stop at a *nonhalt* control. Therefore universality of periodicity or positive return needs the additional source-family premise that every nonhalting input continues forever. It is not an unconditional consequence of the guarded source interface proved here. The halt-pattern and exact-target conclusions do not need that extra premise.

[write] For the literal source of Part II the extra premise holds on its clean input family (Proposition 37.1), and Corollary 38.1 proves that positive return and periodicity are r.e.-complete there, with least period exactly $2\Theta + 2$.

14 A fully materialized accepting example

The release supplies this entire source table: controls **START** and **HALT**, initial control **START**, terminal control **HALT**, and one branch

$$(\text{START}, c_0, c_1) \mapsto (\text{HALT}, c_0 + 1, c_1) \quad \text{for every } (c_0, c_1) \in \mathbb{N}^2.$$

The start has no predecessor and the image at halt is exactly $c_0 > 0$. With $J = 1$, the parameters are

$$m = 2, \quad p = 1, \quad a = 0, \quad D = 8, \quad S = 18, \quad B_2 = 9, \quad L = 28, \quad B_3 = 37, \quad Z = 382.$$

The eight signed gaps, in order, are

$$H_{\text{START}}^{\pm} : 1, 2; \quad H_{\text{HALT}}^{\pm} : 3, 4; \quad O_e^{\pm} : 5, 6; \quad I_e^{\pm} : 7, 8.$$

For input $(0, 0)$, the raw initial and first halt supports are exactly

$$\{-382, 0, 18, 19, 382\}, \quad \{-383, 0, 18, 21, 382\}.$$

There are 47 edge factors and 48 phase factors, for 95 total; the composition radius bound is 21150 and the observer length is 27. Its one source edge takes

$$1 + 346 + 1 + 347 + 1 = 696$$

CA steps: dispatch, outbound travel, endpoint interaction, return travel, commit. The reflection period is 1394. The observer occurs just at time 696 in that cycle. The raw trace includes all occupied coordinates at every time, and the next state after the complete cycle is checked against the initial support, using the actual full gate composition.

The four explicit free/contact demonstrations cover both directions and both boundary cases. In a departure case a triple gate acts at predecessor distance L and its reverse pair competitor is blocked at the same predecessor key. In an arrival case a pair gate acts at predecessor distance $L + 1$, with no contact gate at that predecessor distance. These finite demonstrations support rather than replace the general argument in Section 9.


15 Verification, reproducibility and limitations

The portable archive separates the authored article, finite compiler and source fixtures, certificate implementation and examples, vector renderer, and the companion lower-bound article. It requires only Python's standard library for mathematical replay. Rebuilding the articles additionally requires a normal L^AT_EX installation with the packages declared in their sources. No third-party


A complete reversible five-particle accepting orbit

Each mark is an occupied site from the exact raw trace; blue + phase, red - phase


Full cycle: first halt at 696; return after 1394 steps



Dispatch detail



Halt and sign-reflection detail



Data: sample-orbit.json | exactly five ones in each frame | no particle labels or extra track

Figure 1: Raw five-particle space–time trace of the fully materialized one-increment source. Every mark is drawn from the recorded occupied coordinates; it is not an illustration inferred only from source-level control. The large panel shows the full reflection cycle, and the detail panels expose the close pair near dispatch and halt, where its two particles would otherwise visually merge. The exact trace and deterministic vector renderer are included.

primary-source PDF, screenshot, private working note, execution log, bytecode cache or external dependency installation is bundled.

[write] The companion lower-bound article (Report 12) is not shipped with this report: it is printed in the collection as source 17 of `signal-machine-collision-certificates`. The replay runner and fixtures named in this section are shipped under the prefix `15-rev-five-` (README).

The replay runner checks the SHA256 manifest before testing, stages mutable outputs in a temporary directory, runs the exact checks, and compares reproduced mathematical outputs with frozen fixtures. Its purpose is to make the release independently executable after fresh extraction, with no absolute workspace paths or network dependency. Generated elapsed-time measurements are not mathematical fixtures and are not compared byte for byte.

The final independent CA executable audit comprises ten suites in ordinary Python and again with optimization enabled. It includes 27412 systematic source/image predicate checks, 4448 home-semantics checks, 1080 independent isolated-gate window inputs, and 835 malformed finite inputs tested for both inverse identities and numerical conservation. It also tests strict schema validation, deep and cyclic Boolean syntax, constructor bypasses, immutable snapshots, signed coordinates and exact integer domains. Runtime validation is not implemented by assertions that disappear under optimization.

The independent certificate audit has six suites in each mode. It compares 384 certificate variants, 7008 exact residual rows and 43318 expanded coefficients against separately built formulas; enumerates 6274 bounded natural/rational complete witness tuples; and rejects 556 single-coordinate mutations of accepted witnesses. A fractional convex mixture demonstrates why the extra real selector-norm residual is needed. The producer's separate suite performs 13640 exact checks in each mode. A reinitialization loophole in an earlier supposedly immutable certificate record was found and fixed; the released constructors reject a second initialization before mutation, and both audits passed the final bytes. The formulas themselves were unaffected.

Both accepting raw orbits are regenerated by full CA steps and compared with the recorded finite fixtures. Separate independent scripts reconstruct all 1394 ordinary and 5670 clean-target frames directly from the mathematics. The certificate CLI regenerates natural and paid-real ordinary and clean artifacts and compares each polynomial and witness byte for byte. The vector renderer regenerates its PDF and SVG from the raw ordinary trace.

The independent isolated-swap proof establishes reversibility on the infinite full shift. A periodic-ring exhaustive check corroborates it for one tiny instance: all $2^{17} = 131072$ binary inputs with support-radius parameter $B = 1$ were tested, and 3060 change. Every case preserves raw keys, eligibility, mass and the square-identity law. This is finite corroboration, not an exhaustive check of the general theorem.

The independent geometric compiler stress test uses $m = 5, p = 4, a = 1, J = 2$ and 954 gates. It checks 31957 source-path microedges, 63914 signed endpoint checks and 250 signed homes. Its deliberately merging target control has disjoint *actual* image predicates; it also tests decrement on both sides, stuck states, image gaps, interaction endpoints and the free/contact boundary. These checks isolate potential geometric and predicate errors, but neither this finite source nor the accepting example proves universal source existence.

The release's claims have four distinct supports: the all-input mathematical lemma, the admissible simulation proof, the classical universal-source theorem, and reproducible finite implementation checks. None is silently used in place of another. The software is an audited research implementation, not a certified proof assistant extraction. The mathematical audit does not certify a literal universal source table that is absent from the release.

16 The five-particle threshold and relation to earlier work

The companion Report 12 proves a separate four-mass decision theorem for fixed one-dimensional finite-alphabet conservative CAs with a unique zero-weight vacuum, positive integer nonvacuum weights, and at most one weight-one symbol. For mass at most four, exact finite-configuration reachability and finite-pattern occurrence, anchored or anywhere, are decidable. Its proof handles translating packets and expanding shuttles rather than assuming timed semilinearity. It is included with its original L^AT_EX, PDF and vector figure byte-preserved [10].

Corollary 16.1 (Matched binary reversible particle threshold). *Taking the separate four-mass decision theorem as proved in the companion, five is the least number of ones allowing undecidable effective finite-input halt-pattern occurrence or exact finite-target reachability for a fixed one-dimensional globally reversible, number-conserving binary CA of arbitrary finite radius. For exact reachability, both the initial and target configurations are computable functions of the input.*

Proof. The companion applies to all binary conservative CAs and hence to their reversible subclass, excluding a computable halting reduction at mass at most four. Corollary 4.2 gives such a reduction with exactly five ones. Corollary 13.1 supplies the five-particle exact-target upper bound, and the same companion supplies its mass-at-most-four lower bound. All statements use finite zero-background inputs and their specified local observations or exact finite targets. $\square$

This matched statement explicitly depends on two proofs. The upper construction does not reprove the four-mass theorem. The companion’s numerical four-mass statement is broader in alphabet and does not by itself provide the binary upper bound. Its byte-preserved version predates the present reversible binary upper construction, so its discussion of that previously open part is read in that chronological context.

[write] Report 12 is source 17 of `signal-machine-collision-certificates` (Part VI, Theorem `smc:fm:thm:main`); it arrived in batch 80, a day before this report. Its statement that a binary five-particle upper bound was open is answered by this Part: Report 14 (second half of this Part) and the present Report 15 each supply it.

Morita–Imai’s 2001 construction provides a finite-zero-background, number-conserving reversible simulation with a numerical triple-part alphabet; its reported mass depends on the source-control count [3]. Morita’s 2012 construction gives a one-dimensional reversible number-conserving CA with 96 numerical states, but its discussed universal configurations have an infinite ultimately periodic background [4]. Neither statement alone is a binary five-one construction. These are relevant precedents, not inputs to the isolated binary-swap proof.

The 2016 article by Alhazov and Imai is directly relevant to particle complexity [5]. Its full construction was not available in the targeted source review used here, so no inference about its exact binary, reversible or optimality scope is made. An unsuccessful or incomplete search is not a novelty or priority argument. No claim of being the first such construction is made in this report.

17 Research questions and sharpenings

The construction raises concrete questions that do not affect the established bounds:

1. **Literal universal source.** Materialize and independently audit a universal reversible two-counter table in the accepted schema. This would turn symbolic m, p, a, J into numerical universal parameters and permit end-to-end source-level replay of a chosen encoded computation. [write] Answered by Part II (Report 16): a literal universal reversible

table in this schema, with $m = 122,622$, $p + a = 141,561$ and $J = 0$, its clean loader and all-input proofs. End-to-end source-level replay of a universal computation remains out of reach (Report 16 predicts about $8 \cdot 10^{13}$ source steps for empty-tape initialization; Report 17 reduces that startup to 138 executed steps).

2. **Radius and gate optimization.** The sum-of-factor bound is intentionally generous. Can one reuse gap codes, reduce isolation halos, or prove a smaller composed dependency cone without losing all-input reversibility? Can the fixed five-particle construction be made near a useful minimal radius?
3. **More economical certificates.** Finite class normalization can multiply branch count by $(J + 2)^2$. Can a direct guard-circuit encoding preserve a unique natural witness fiber with fewer residuals and controlled degree?
4. **Uniform versus bounded arithmetization.** The horizon here is compiled syntax. A single horizon-free polynomial representing the same unbounded reachability relation would require an additional genuine uniform arithmetization theorem, not a reinterpretation of these bounded certificates.
5. **Mechanized proof.** Formalize the all-input raw-key invariance lemma, enabled-edge admissibility, and signed matching construction in a proof assistant; relate the executable sparse-set implementation to its local-rule semantics.
6. **Conservation and small-particle lower bounds.** Examine how the separate four-mass analysis changes with multiple distinguishable unit labels, nonuniform schedules, higher-dimensional lattices or active zero-weight backgrounds. The current threshold does not cover those different models.

Five particles in a binary conservative cellular automaton

An explicit counter compiler and finite horizon Diophantine certificates

3 October 2026

Abstract of Report 14, as delivered. Every finite deterministic two-counter program with increment and conditional decrement compiles effectively to a fully specified one-dimensional binary conservative cellular automaton. Every valid simulation uses exactly five occupied sites. Finite geometric gap codes carry control; two unbounded gaps carry counters. A read-only remote zero sensor and strictly bounded component writes make the rule conservative on arbitrary binary configurations, including malformed inputs. The proof supplies a finite radius, exact instruction clocks, and a finite anchored halt word with no premature occurrences.

A pinned 8,408-instruction universal source gives alphabet size two, radius 756,787 and halt-word length 50,452. Its guarded component description has 1,272,679,260 indexed cases. The binary truth table has $2^{1,513,575}$ input rows and is specified algorithmically, not materialized. Universality uses the independently checked finite-input Neary–Woods source and all-input Turing-machine-to-counter proofs, rather than an inference from table size. Together with the separately proved mass-four decision theorem in companion Report 12, this gives a sharp five-particle threshold in the stated binary finite-input halt-observation model.

For each syntactically fixed literal-source horizon h , an explicit ordinary sum-of-squares polynomial has $10,750h$ natural witness variables, $2,344h + 1$ squared residuals and degree at most four. Its complete witness fiber is empty or a singleton. Exact CA time costs one further variable and square; a paid nonnegative-real variant costs one further square per source step. Natural external inputs, the prime-encoded loader, enormous geometric constants and expansion costs are all retained. This is a mathematical proof with executable audits, not a proof-assistant formalization or a priority claim.

[write] Delivered as `five-particle-binary-portable.zip` (batch 82, archive 23; arrival `db37d18c8`), dated 3 October 2026; files placed by `bca6383e9` with the prefix given in Appendix A. Report 14 is unnumbered in its own header; Report 16’s compiler proof identifies it (“Report 14’s nonreversible 8408-row source”). It and Report 15 were written about forty minutes apart and neither cites the other; Report 15’s Corollary 16.1 implies Report 14’s Corollary 27.1 (note after it). Both are printed in full, since their compilers, classes and universal instances differ.

18 The model and the result

A binary configuration is a function $x : \mathbb{Z} \rightarrow \{0, 1\}$. Its occupied set is $X = \{i : x(i) = 1\}$. When X is finite, its mass is $|X|$; thus a particle here is exactly one occupied lattice site, without a hidden color, weight or label. A radius- R cellular automaton (CA) is determined by a Boolean function $f : \{0, 1\}^{2R+1} \rightarrow \{0, 1\}$, with $F(x)(i) = f(x(i-R), \dots, x(i+R))$. Its radius need not be minimal. We require the all-zero vacuum to be fixed and mass to be conserved on every finite-support input. We also give a full-shift transport interpretation, where total cardinality need not be finite.

A source program has a finite set Q of controls, an entry control q_0 , a distinguished halt control q_* with no outgoing row, and two counters $a, b \in \mathbb{N} = \{0, 1, 2, \dots\}$. Every other control has one row:

$$\begin{aligned} \text{ADD}(j, q') &: c_j \leftarrow c_j + 1, \quad q \leftarrow q', \\ \text{SUB}(j, q_+, q_z) &: c_j > 0 : c_j \leftarrow c_j - 1, \quad q \leftarrow q_+; \quad c_j = 0 : q \leftarrow q_z, \\ \text{NOP}(q') &: q \leftarrow q'. \end{aligned}$$

Here $c_0 = a, c_1 = b$ and $j \in \{0, 1\}$. The optional no-operation row is included for convenience. All destinations are defined controls; self-loops are allowed. The implementation orders control

labels lexicographically, with the halt control placed last. Control labels are finite strings, and the supplied source table is held fixed after compilation.

Theorem 18.1 (Binary five-particle compiler). *From every such finite source program one can effectively construct a fixed binary finite-radius CA F , a computable finite-input encoding $\text{Enc}(q, a, b)$ of mass exactly five, and a fixed finite binary word w_* anchored at the origin, with the following properties.*

1. *F is defined on the whole binary shift and conserves every finite-support configuration. Every valid simulation has exactly five particles at every tick.*
2. *There is an exact positive integer duration for each nonhalting source step. The CA reaches the encoding of its source successor after that duration. Iterating gives exact simulation of every finite or infinite source run.*
3. *On every valid encoded orbit, w_* occurs at its fixed anchor precisely when a committed source halt configuration has been reached. The first occurrence time is the sum of the exact source-step durations. Halt persists.*

The theorem states a simulation of sequential source computations. It does not assert reversibility, intrinsic universality, efficient input preparation, efficient CA time, or minimal radius. The proof is independent of whether the chosen source program is universal. Section 26 supplies a particular universal instance with a separately checked provenance chain. Section 27 then combines this upper bound with a separate lower bound in an exactly matched model.

19 Finite codes and buffered geometry

Let $m = |Q|$, let p be the number of ADD or SUB rows, let n be the number of NOP rows, and let z be the number of SUB rows. Thus $m = p + n + 1$. Make one home code H_q for each control and one outbound code O_q and inbound code I_q for each moving control. Give these codes distinct integer gaps, first all home codes, then all outbound codes, then all inbound codes, in the specified control order. Write $d(\cdot)$ for this bijection onto $\{1, \dots, D\}$, where

$$D = m + 2p, \quad C = D + 1, \quad E = 2D + 2, \quad (19.1)$$

$$K = 4D + 5 = 2E + 1, \quad Z = 4K, \quad R = Z + 3K + E = 30D + 37. \quad (19.2)$$

The constants have distinct functions. D bounds a packet's internal gap; C separates it from a marker; E bounds the write enlargement; K defines components and contact; Z is the counter buffer and zero-read distance; R is the sufficient local radius.

For a marker at o , a side $s \in \{-1, +1\}$ and a gap code d , define

$$\mathcal{P}(o, s, d) = \{o + sC, o + s(C + d)\}.$$

This is a two-particle packet. Its two particles have no separate labels. The valid source-boundary encoding is

$$\text{Enc}(q, a, b) = \{-(Z + a), 0, Z + b\} \cup \mathcal{P}(0, +1, d(H_q)). \quad (19.3)$$

The outer left marker encodes a , the outer right marker encodes b , and 0 is the origin marker. These names describe an invariant of valid runs; they are not extra states available to the binary local rule.

All five coordinates in (19.3) are distinct: the close packet is strictly right of 0, its furthest site is at most $C + D = 2D + 1 < Z$, and both outer markers lie at distance at least Z from 0. The only adjacent occupied gap at most D is the packet gap. Every unspecified site is exactly zero.

Translation of the encoding simply translates the whole future and its observer anchor. For definiteness, inputs use origin 0.

If q acts on counter j , put $s_q = -1$ when $j = 0$ and $s_q = +1$ when $j = 1$. The travel direction of O_q is s_q , and that of I_q is $-s_q$. Direction is recovered from a geometric code and the fixed source table; it is not stored as a separate particle property.

20 The complete rule on arbitrary configurations

Join consecutive occupied sites when their distance is at most K . The resulting maximal connected occupied sets are called components. For an infinite configuration, components may be finite, infinite in one direction, or infinite in both directions. Empty regions contain no component. Every component is left fixed unless one of the rules below applies. In particular, singletons, components of size at least four, and infinite components are all frozen.

A three-particle component is eligible only if exactly one of its two adjacent gaps is at most D and the other adjacent gap is in $[C, K]$. The close pair is the packet and the remaining particle is the marker. This parsing is unique because $C = D + 1 > D$. Denote the marker by o , the packet by P , its internal gap by d , and its nearest distance to the marker by $a = \min_{x \in P} |x - o|$. In particular, no triple with two close gaps is parsed as two competing packets.

20.1 Free travel and departure

For a two-particle component with gap $d(O_q)$ or $d(I_q)$, translate both particles by one site in the recovered travel direction. A home-coded pair is fixed. Any pair without a recognized travel code is fixed.

In an eligible travel triple, let $v \in \{-1, +1\}$ be its travel direction. The marker is *behind* the packet when

$$(v = +1 \text{ and } o < \min P) \quad \text{or} \quad (v = -1 \text{ and } o > \max P).$$

In that case fix the marker and translate the packet by v , for every $a \in [C, K]$. If the marker is ahead, do nothing unless $a = K$; the contact rules below cover exactly that case. An ahead marker at distance less than K freezes. Such malformed cases need not simulate anything, but have a well-defined conservative evolution.

20.2 Home dispatch and the contextual zero test

A home triple must have the exact orientation and coordinates

$$\{o, o + C, o + C + d(H_q)\}.$$

Other arrangements carrying a home gap freeze. The home rules retain o :

1. For $\text{ADD}(j, q')$, replace the packet by $\mathcal{P}(o, s_q, d(O_q))$.
2. For $\text{SUB}(j, q_+, q_-)$, inspect the old binary value at $o + s_q Z$. If that value is one, replace the packet by $\mathcal{P}(o, +1, d(H_{q_+}))$. If it is zero, dispatch $\mathcal{P}(o, s_q, d(O_q))$.
3. For $\text{NOP}(q')$, replace it by $\mathcal{P}(o, +1, d(H_{q'}))$.
4. For $q = q_*$, leave the component fixed.

The contextual inspection is a read of the old global configuration. It does not move or erase the sensed site, does not reserve its component, and does not change the component partition. All components are updated simultaneously from the same old configuration.

20.3 Endpoint reversal and home commit

Consider an ahead-marker contact at distance exactly K . For code O_q , write $s = s_q$ and

$$\Delta_q = \begin{cases} +1, & q \text{ is ADD,} \\ -1, & q \text{ is SUB.} \end{cases}$$

Move the contacted marker to $o' = o + s\Delta_q$, and replace the packet by $\mathcal{P}(o', -s, d(I_q))$. This reverses the packet on the inner side of the updated endpoint. On valid runs the decrement case was dispatched only when the selected counter was positive.

For code I_q , retain the contacted marker o and put the packet at $\mathcal{P}(o, +1, d(H_{q+}))$, where q_+ denotes the ADD successor or the positive SUB successor. In a left-counter excursion this final rewrite may carry the packet across the origin marker. The marker itself remains fixed. All unlisted cases freeze.

Lemma 20.1 (Deterministic parsing). *These rules assign exactly one output set to every component in every binary configuration, without a promise of valid encoding.*

Proof. Component sizes separate pair rules from triple rules. An eligible triple has one uniquely selected close pair and one marker. Its positive gap selects exactly one code among the disjoint H , O , and I code ranges. Home geometry has a specified orientation and distance. Travel direction is fixed by the code; ahead and behind are exclusive. A contact requires precisely $a = K$. Finally, the two home SUB alternatives read the same old bit and have complementary guards. No case requires a priority ordering between competing rewrites. Every remaining configuration is covered by the freeze convention. $\square$

21 Conservation with remote reads and bounded writes

The zero sensor can lie in a different component, whose own update can be unrelated to the computation. This is the reason to separate reading distance from writing distance. Conservation depends on a uniform bound on the latter, not on independence of the former.

Lemma 21.1 (Equal cardinality and bounded output hull). *For every finite component A , its replacement A' has $|A'| = |A|$ and*

$$A' \subseteq [\min A - E, \max A + E].$$

The statement holds for either value of every contextual guard.

Proof. Frozen components retain all their sites. A translated pair retains two distinct sites. Every rewritten triple consists of a marker and a pair with positive internal gap and marker-to-pair distance at least $C > D$, hence has exactly three distinct sites.

Free and departure translations enlarge the hull by at most one. A changed home gap enlarges it by at most D . A dispatch to the opposite side or a return commit across the marker has new coordinates of absolute marker relative value at most $C + D = 2D + 1 < E$.

At an outbound contact the old near packet site lies K on the inner side of the marker, and the far site is another $d(O_q)$ inward. The new marker moves by one, while both new packet sites are on its inner side within $C + D$. The inequalities $K > C + D$ and $K > C + 2D + 1$ put the new inner pair between the old far extent and the marker, allowing the marker's one-site enlargement. Thus the coarser common allowance E covers every case. Both zero-test outcomes satisfy the same estimate, regardless of what caused the read bit. $\square$

Proposition 21.2 (Global numerical conservation). *The component rule conserves mass on every finite-support binary configuration.*

Proof. Two distinct old components have occupied endpoint separation greater than K . Their closed output-hull allowances enlarge toward each other by at most $2E$, whereas $K = 2E + 1$. The enlarged lattice intervals are therefore disjoint. This includes intervals around frozen components. Consequently output sites from different components cannot collide or overwrite retained sites. Summing the cardinality equality of Lemma 21.1 over components gives $|F(X)| = |X|$. No valid-input invariant entered the proof. $\square$

For example, two malformed home triples may simultaneously read the same remote particle. Those reads may all select a misleading zero branch. Each selected rewrite still preserves its own cardinality and stays in its private bounded output allowance. Shared reads do not consume the particle, so they cannot create or destroy mass. Dense components and ambiguous packets are frozen instead of being resolved by an unproved global scheduler.

There is also a useful full-shift formulation. Match old and new particles in increasing order within each changed component; match frozen particles to themselves. Each changed component has diameter at most $2K$, so its matching has displacement at most $2K + E$. Disjoint enlarged hulls make the union a bijection from the old occupied set to the new occupied set, for arbitrary infinite as well as finite configurations. This matching is translation equivariant. On a periodic configuration it descends to a bijection of occupied site orbits modulo the spatial period, proving per-period conservation. This is a transport statement, not an assertion that the CA map itself is injective or reversible.

22 A finite local rule and its radius

A component definition is not automatically a local rule: in particular, the instruction to freeze an infinite or very large component must be locally indistinguishable from the appropriate finite cases. We verify that the claimed radius covers both the positive detection and this exclusion issue.

Lemma 22.1 (Local determination). *The output bit at i is determined by the old word on $[i - R, i + R]$, with $R = Z + 3K + E$ from (19.2).*

Proof. A component that can change or supply the output at i has at most three particles. By Lemma 21.1, its old hull is within E of i , and its diameter is at most $2K$. Thus all its particles are within $E + 2K$ of i . To certify that it is an entire component, inspect the K sites just beyond each endpoint. These boundary guards are within $E + 3K$ of i . Any contextual home read is at distance at most Z from its marker, hence within $E + 2K + Z$ of i . Both bounds are at most R .

Conversely, a large component containing a relevant occupied site cannot be misread as a small one merely because the observed word is truncated. If it contains at least four particles, four consecutively connected particles including that site can be exposed within distance $3K$ of it. Equivalently, if a purported size-two or size-three component of the truncated word could change the central output, the preceding bounds put *both* its complete K -wide boundary guards strictly inside the observation window. It therefore has no hidden continuation outside that window. Components whose only observed fragments lie near the window endpoints cannot write at i , by the same diameter and hull allowance. For a central occupied site in a frozen large component, the four-site witness certifies that it stays occupied. An empty central site can acquire a particle only through one of the completely visible small components. This handles retention as well as creation of the central bit and proves equality with the full-configuration definition. $\square$

A literal Boolean evaluator is now specified: given the exact radius- R word, regard its ones as a finite set of occupied offsets, partition them into K -components, apply Section 20 using the supplied old word, and return one exactly when its output contains offset zero. The code `BinaryCA.local` implements this finite algorithm. Absent offsets in the supplied window mean exact zeros; offsets outside the window are rejected. Lemma 22.1 proves that this evaluator agrees with the full-shift map. Translation equivariance and time independence follow from the rule's relative-coordinate definition. The vacuum is fixed.

The count of Boolean inputs is exactly 2^{2R+1} . A finite program that terminates on every one of those inputs completely specifies the local table. This release does not enumerate that bit string. The sparse component evaluator is a convenient way to compute selected local outputs and whole finite traces; it is not a different state space or an extra oracle outside the CA.

23 The all-input simulation proof

We now restrict to valid encodings. The source counters may be arbitrarily large natural numbers; none of the proof depends on a finite test range.

Lemma 23.1 (Packet and marker separation). *Throughout a valid source excursion, the packet has internal gap at most D , there is exactly one close pair, and it cannot belong to a component containing two distinct markers. All other components are isolated markers.*

Proof. At a source boundary the markers are at $-(Z+a)$, 0 , $Z+b$ and their consecutive separations are at least Z . Home and travel codes all have gaps in $[1, D]$. A two-site packet can connect two markers by gaps at most K only if their separation is at most $2K + D$. But $Z = 4K > 2K + D$. Endpoint addition increases its separation from the origin; positive decrement reduces it to no less than Z . The origin stays fixed, and both reversals and home commits place an intact coded packet at distance at least $C > D$ from their marker. Translations preserve its internal gap. These facts propagate the asserted separations through every phase. $\square$

Lemma 23.2 (Exact home zero test). *At a valid home configuration with selected counter c and side s , the old bit at sZ is one if and only if $c = 0$.*

Proof. The selected endpoint is at $s(Z + c)$, so it occupies sZ precisely when $c = 0$. The other endpoint is on the opposite side, the origin is at zero, and the entire home packet is within $C + D < Z$ of zero. No other valid particle can occupy the tested site. Thus dispatch of a decrement excursion implies $c \geq 1$; it can never decrement below the buffer. $\square$

23.1 Exact microstates of an arithmetic excursion

Fix a moving control q , selected old counter c , side $s = s_q$, and change $\Delta = +1$ for addition or -1 for positive decrement. Put

$$d_O = d(O_q), \quad d_I = d(I_q), \quad (23.1)$$

$$L_O = Z + c - K - C - d_O, \quad L_I = Z + c + \Delta - K - C - d_I. \quad (23.2)$$

The untouched endpoint stays fixed throughout. After the dispatch tick the packet is $\mathcal{P}(0, s, d_O)$. After u subsequent travel ticks, $0 \leq u \leq L_O$, it is

$$\{s(C + u), s(C + d_O + u)\}. \quad (23.3)$$

Its separation from the origin initially equals C and grows one site at a time. It first uses the behind-marker rule, then the free-packet rule. At $u = L_O$ its leading site is exactly K from the

selected endpoint $s(Z + c)$. Integer unit motion cannot skip this contact. Before then its ahead distance is greater than K , so there is no early reversal.

The next tick moves the endpoint to $o' = s(Z + c + \Delta)$ and places the inbound packet at $\mathcal{P}(o', -s, d_I)$. After v return travel ticks, $0 \leq v \leq L_I$, its sites are

$$\{o' - s(C + v), o' - s(C + d_I + v)\}. \quad (23.4)$$

Again it departs a behind marker, travels freely, and then contacts the origin at distance K , exactly at $v = L_I$. Its next tick commits the successor home code. The resulting configuration is $\text{Enc}(q_+, a + \Delta, b)$ for the left counter and $\text{Enc}(q_+, a, b + \Delta)$ for the right counter.

The travel lengths are strictly positive even at the smallest counters:

$$L_O \geq Z - K - C - D = 3K - C - D > 0,$$

and, using $c + \Delta \geq 0$ in the decrement case, $L_I \geq 3K - C - D > 0$. Thus the formulas include departure, separation and contact without a zero-length ambiguity. At every intermediate phase the packet lies between the two outer markers. The origin-crossing commit also remains there.

Proof of Theorem 18.1 apart from observation. The whole-shift CA and conservation assertions follow from Proposition 21.2 and Lemma 22.1. At a home state, Lemma 23.2 selects the correct zero branch or arithmetic excursion. The explicit microstates (23.3)–(23.4) and Lemma 23.1 show that all intermediate rules apply exactly as claimed and that the endpoint update changes only the selected counter. A zero decrement or no-operation instead commits its stated successor in one tick. A halt encoding is fixed. Induction over source steps proves the source-boundary invariant for arbitrary natural inputs. Because every nonhalting source instruction has a finite positive duration, any infinite source run produces an infinite CA run with all its source boundaries reached. Global conservation keeps mass exactly five at every intermediate tick. $\square$

24 Exact clocks and the anchored halt word

An arithmetic excursion consists of one dispatch, L_O translations, one endpoint reversal, L_I translations and one home commit. Therefore its exact duration is

$$\tau_q(c) = 3 + 2(Z + c) + \Delta_q - 2K - 2C - d(O_q) - d(I_q) = 2c + \kappa_q, \quad (24.1)$$

where $\kappa_q = 3 + 2Z + \Delta_q - 2K - 2C - d(O_q) - d(I_q)$. Zero decrements and no-operations take exactly one tick. A state already halted has first-halt time zero, and its future configuration is fixed.

Define a finite word w_* on the integer interval $[0, C + D]$ by requiring ones exactly at

$$0, \quad C, \quad C + d(H_{q_*}), \quad (24.2)$$

and zeros at *every* other site of the interval. Its length is $C + D + 1 = 2D + 2 = E$. No counter value appears in this word.

Proposition 24.1 (No false or early halt observation). *On the orbit of any valid encoding, the anchored word w_* occurs if and only if a committed source halt has been reached. Its first occurrence is at the sum of the durations (24.1) and the unit zero/no-operation durations along the first-halt source trace.*

Proof. Both endpoints remain outside the observer window. The three prescribed ones must therefore be the origin and the packet. Their only close pair has gap $d(H_{q_*})$. At every intermediate

time there is exactly one close pair, and every outbound and inbound code has a different gap. The halt gap also differs from every other home code. Hence an excursion whose *future* source successor is halt cannot exhibit the word before the home commit. At a committed home state the packet has the exact positive-side geometry of (24.2), and the word appears precisely for control q_* . The required zeros hold because no additional particle lies in the window. The halt rule is the identity, so the word then persists. Summing the exact step clocks gives its first occurrence time. $\square$

This completes Theorem 18.1. The observation is anchored and finite. It is not an exact finite-configuration target test: the two final counter endpoints can be arbitrarily far away and can encode arbitrary final counter values. Conversely, conservation does not weaken the observer's zero requirements; they are essential parts of the finite word.

25 Exact finite description counts

Before removing identity cases or cases unreachable from valid encodings, the complete component scheme has the following indexed guarded cases:

Rule family	Number of indexed cases
Free outbound and inbound packet translation	$2p$
Behind-marker departure for every $a \in [C, K]$	$2p(K - C + 1)$
Ahead-marker endpoint contact and home contact	$2p$
Home dispatch or commit, splitting each zero read	$p + n + z$
Total	$4p + 2p(K - C + 1) + p + n + z$

Every SUB has two complementary home guard alternatives, which accounts for the extra z beyond the number of nonhalting controls. Each travel code has one permitted behind orientation and one ahead orientation, so no uncharged direction label is missing. The otherwise-frozen convention is part of the finite rule definition.

These are counts of an indexed component schema. This release stores the finite algorithm and its source table, not one written record for every indexed component case. In particular, neither this schema count nor the number of source instructions is the number of Boolean local-table rows. The latter is exactly 2^{2R+1} . Given a neighborhood, the algorithm fixes its central output in finite time, so the distinction is representational rather than a missing rule or unspecified behavior.

26 The pinned universal source

[write] **A second presentation.** The 8,408-row source of this section is the prime-coded program of quadratic-orthant-certificates (QOC) Part V, source 16 (its `$qoc:um:sec:virtual` and `$qoc:um:sec:prime`); the 528-row three-counter program, the stack invariants with cost $5Q + r + 7Y_0 + w + 4$, the prime macros with cost $(3p + 1)N + 2$, the row count 8,408 and the verifier's 4,475 symbolic paths are proved there as Lemmas `qoc:um:lem:virtual` and `qoc:um:lem:prime`. Report 14 shipped that report's `verify_source.py` and receipt byte for byte; they are not shipped again. The section is kept in full, as Report 14's own route to the same lemmas, because Report 14's universal instance and its counts depend on it; no novelty is claimed for it.

The generic theorem becomes a universal CA only after a universal source is supplied and its input interface is matched to the encoding. The delivered source has 8,408 literal two-counter instructions, comprising 6,068 ADD rows and 2,340 SUB rows, with 8,409 controls including halt. All aliases and instruction macros have been expanded. Its exact bytes are pinned in the portable source-verification dependency. Row count alone establishes none of its semantic or universality properties.

26.1 The finite-input primary machine

The primary reference is T. Neary and D. Woods, *Four Small Universal Turing Machines, Fundamenta Informaticae* **91** (2009), 105–126, DOI 10.3233/FI-2009-0008 [1]. Equation (3), Definition 3.1 and Table 1 give effective finite program/data encodings with blank tails, start control u_1 , and scanned blank c . Section 3.5 and Table 16 supply the corresponding 15-state, two-symbol machine. This use is ordinary finite-input universality; it does not require an infinite periodic background. Both finite half tapes may vary with the encoded program and input.

The literal table uses $c = 0$, $b = 1$ and $u_1, \dots, u_{15} = A, \dots, O$. For auditability, Table 5 lists all 30 entries. Each defined entry is (written symbol, direction, next control); undefined $J1$ is the halt. The independently checked serialization is bundled as `tm_table.json`.

Table 5: The pinned source TM, matching primary Table 16

Control	Read 0	Read 1	Control	Read 0	Read 1
A	$0, R, B$	$1, R, A$	I	$0, R, A$	$1, L, J$
B	$1, R, C$	$1, R, A$	J	$1, L, K$	undefined
C	$0, L, G$	$0, L, E$	K	$0, R, L$	$1, R, N$
D	$0, L, F$	$1, L, E$	L	$0, R, M$	$1, R, L$
E	$1, R, A$	$1, L, D$	M	$0, L, B$	$1, R, L$
F	$1, L, D$	$1, L, D$	N	$0, L, C$	$0, R, O$
G	$0, L, H$	$1, L, G$	O	$0, R, N$	$1, R, N$
H	$1, L, I$	$1, L, G$			

The primary scan contains a discrepancy that must remain visible. Table 16, printed page 121 (PDF page 17), leaves u_{10}, b undefined. The displayed halting configuration on printed page 123 (PDF page 19) also underlines b . The prose immediately below instead says u_{10}, c . The table and display agree with the pinned $J1$ halt used here; the prose is recorded as inconsistent. No author-issued erratum is asserted, and no repair to the serialized table has been inferred. The source’s old bibliography comment also gives the incorrect page range 123–144. That comment is preserved to retain the pinned bytes, while the report and provenance file give the verified 105–126 range. No third-party paper or screenshot is redistributed in this package.

[write] **Correction (3 October 2026)**. The bibliographic details in this subsection are wrong, and the “old bibliography comment” is right: the publisher’s record of Neary–Woods is *Fundamenta Informaticae* 91(1) (2009), 123–144, DOI 10.3233/FI-2009-0036, while DOI 10.3233/FI-2009-0008 is a different paper and 105–126 is the author-hosted PDF’s own pagination (Appendix A, entry [1]). The page locators “printed page 121” and “printed page 123” refer to that PDF. The (u_{10}, b) discrepancy is the one recorded in quadratic-orthant-certificates after its Table `qoc:wf:tab:tm`, and Table 5 agrees with that table entry by entry.

26.2 Finite half tapes and three natural counters

Let L, R store the finite bit stacks strictly left and right of the head, least significant bit nearest the head; omitted high bits are blank zero. The scanned bit is finite control, and a third register T is scratch. At a genuine TM cut, $T = 0$. The literal virtual source starts at a prologue that clears arbitrary T with

$$\text{SUB}(T, \text{clear}, \text{first TM control}).$$

Repeated positive branches remove one unit, followed by one failed test. This takes $T + 1$ virtual instructions and enters control $A0$.

For a defined TM transition $(q, s) \mapsto (w, D, q')$, denote the half tape in direction D by X and the other half by Y . Write $X = 2Q + r$, $r \in \{0, 1\}$. The exact stack action is

$$X \leftarrow Q, \quad Y \leftarrow 2Y + w, \quad T \leftarrow 0, \quad \text{new finite control } (q', r). \quad (26.1)$$

The 528-row virtual three-counter table implements this using the following finite ADD/SUB graphs. Their labels are distinct for each source transition and parity branch; the delivered files contain the fully expanded rows.

The pop loop tests and decrements X twice, then adds one to T and repeats. A failed first test selects remainder 0; a successful first test followed by a failed second test selects remainder 1. After k complete pair iterations,

$$X = 2(Q - k) + r, \quad T = k.$$

The rank $Q - k$ strictly decreases on full iterations. At exit $X = 0, T = Q$, with r in control. A transfer loop repeatedly decrements T and increments X until $T = 0$, producing $X = Q$. The pop stage uses exactly $3Q + (1 + r) + (2Q + 1) = 5Q + r + 2$ virtual instructions.

To push the written symbol, drain Y into T at two increments of T per successful decrement of Y . After k iterations, $Y = Y_0 - k, T = 2k$, with decreasing rank $Y_0 - k$. Next transfer T back to Y one unit at a time; its invariant is $Y + T = 2Y_0$, with decreasing rank T . This reaches $Y = 2Y_0, T = 0$. If $w = 1$, one final ADD writes the low bit; if $w = 0$, the graph links directly to the next TM cut. This stage uses $7Y_0 + 2 + w$ virtual instructions. Consequently the full TM transition costs

$$C_{\text{virtual}} = 5Q + r + 7Y_0 + w + 4. \quad (26.2)$$

These invariants prove (26.1) for all natural half tapes. Every loop terminates, and none can halt internally. The undefined $J1$ transition alone reaches HALT. Induction supplies finite-prefix simulation and halting equivalence for the complete TM run.

A repeated control label is not, by itself, a TM cut: the pop loop revisits its first label while $T > 0$. Correct replay tests the cut label *and* $T = 0$. This qualification also matters at the next compilation layer.

26.3 Three counters in two by prime coding

At every virtual-instruction boundary the physical counters satisfy

$$A = C_0 2^L 3^R 5^T > 0, \quad B = 0, \quad \gcd(C_0, 30) = 1. \quad (26.3)$$

Unique factorization gives such a representation for every positive raw input A . The cofactor C_0 is unchanged. The prologue clears the exponent of 5, so the raw input starts the TM with $L = v_2(A)$, $R = v_3(A)$ and scratch zero regardless of its other factors. No external factorization is executed by the loader or by the counter machine. Intermediate physical A may be zero; positivity is required at virtual cuts.

For a virtual ADD at prime $p \in \{2, 3, 5\}$, drain physical $A = N$ one unit at a time, placing exactly p increments into B after each successful SUB of A . After k iterations,

$$A = N - k, \quad B = pk.$$

The rank $N - k$ proves termination at $A = 0, B = pN$. A second loop transfers B one-for-one back to A , with invariant $A + B = pN$ and decreasing rank B . It exits at $A = pN, B = 0$, exactly increasing the chosen exponent. Counting literal rows executed gives $2pN$ ADDs, $(p + 1)N$ successful SUBs and two failed SUBs, for a total

$$C_{\text{inc},p}(N) = (3p + 1)N + 2. \quad (26.4)$$

For a virtual SUB, use p finite remainder controls. Each successful SUB of A advances the remainder cyclically; crossing from $p - 1$ to 0 additionally increments B . At remainder control j after k complete groups,

$$A = N - pk - j, \quad B = k, \quad 0 \leq j < p.$$

Each successful test consumes a unit of A , so this phase terminates. Writing $N = pQ + r$, $0 \leq r < p$, the failed test occurs at $A = 0, B = Q$ in remainder control r .

If $r = 0$, transfer B to A one-for-one and take the positive virtual branch. The result is $A = Q, B = 0$; since $N > 0$ and $p \mid N$, $Q \geq 1$. Divisibility means precisely that the chosen exponent was positive, and division by p decrements it. If $r > 0$, drain B with p ADDs to A per successful SUB, then append r further ADDs to A . After k restoration groups, $A = pk, B = Q - k$, so the result is $A = N, B = 0$. Take the zero branch. Nondivisibility means the exponent is zero, so this branch restores the entire encoded state. Remainders are bounded control, never a third unbounded physical register. All transfer ranks decrease and all continuations restore $B = 0$ before leaving the macro.

The exact SUB costs are

Remainder	ADD	Successful SUB	Failed SUB	Total
$r = 0$	$2Q$	$N + Q$	2	$N + 3Q + 2$
$r > 0$	$N + Q$	$N + Q$	2	$2N + 2Q + 2$

A virtual ADD uses $p + 3$ literal rows. A virtual SUB uses $(3p^2 + p + 4)/2$ rows, namely 9, 17 and 42 at primes 2, 3 and 5. The virtual source has ADD counts (74, 76, 145) and SUB counts (58, 58, 117) on (L, R, T) , including the clearing prologue. Therefore the literal row count is

$$74 \cdot 5 + 76 \cdot 6 + 145 \cdot 8 + 58 \cdot 9 + 58 \cdot 17 + 117 \cdot 42 = 8,408.$$

The actual literal file, rather than an abstract row-count assertion, fixes every intermediate label and destination. A physical macro cut requires its entry label *and* $B = 0$, since internal drain loops can revisit that label while $B > 0$.

26.4 What the source verifier proves

The portable source verifier reads the saved tables and does not import the source generator. It checks affine straight-line loop bodies such as

$$(A, B) = (x + p, y) \longmapsto (x, y + 1)$$

and their exits, transfers and bounded tails. Parameters are arbitrary naturals. Every SUB on a verified path has either identically zero input or an affine value with nonnegative coefficients and

constant at least one. Thus every selected branch and affine endpoint equality holds uniformly, not just at sampled parameter values.

There are 4,475 checked symbolic paths, covering all 528 virtual and 8,408 literal instructions. The decreasing-rank and invariant proofs above compose these paths into arbitrary-length terminating loops. This is the all-input argument. Separate finite checks include 1,536 prime-macro cases, 4,698 TM/virtual cases, 261 whole TM transitions through the literal physical table, and 81 raw-cofactor/scratch-prologue cases. Such tests supplement the proof; they do not certify universality by extrapolation.

26.5 The resulting fixed universal CA

For a finite primary TM encoding, construct its natural stacks L, R and use physical input $(A, B) = (2^L 3^R, 0)$. Then use (19.3) for the fixed literal source entry control. This is an effective finite-support encoding of exactly five particles. Composing primary finite-input universality, the two all-input counter simulations, and Theorem 18.1 gives one fixed binary CA for which anchored halt-word occurrence is undecidable on this computable family of finite inputs.

The exact instance has

Quantity	Value	Quantity	Value
m	8,409	p	8,408
z	2,340	D	25,225
C	25,226	E	50,452
K	100,905	Z	403,620
Radius R	756,787	Observer length	50,452

The component-family counts are 16,816 free cases, 1,272,634,880 departure cases, 16,816 contact cases and 10,748 guarded home cases. Their sum is 1,272,679,260. The Boolean local truth table instead has

$$2^{2R+1} = 2^{1,513,575} \quad \text{input rows.} \quad (26.5)$$

Neither the billion indexed component cases nor this binary truth table is expanded as a flat list in the archive. All of their outputs are fixed by the small rule algorithm and pinned literal source.

The counter input can be extraordinarily large: its binary bit length is $\lfloor L + R \log_2 3 \rfloor + 1$ for $A = 2^L 3^R$, while its geometric left endpoint is $-(403,620 + A)$. L and R themselves are integer values of binary stacks, not merely their bit lengths. Prime-macro instruction costs depend on these numeric values by (26.4) and the SUB ledger; CA clock costs depend on physical counters by (24.1). The construction does not absorb the external loader into a free operation or claim a polynomial-time map in the length of the original TM input.

27 The sharp particle threshold and its scope

Companion Report 12, *Four mass units and exact reachability* [10], proves decidability of finite-pattern occurrence, anchored or anywhere, from finite inputs of mass at most four in a broader finite weighted-alphabet model. Its hypotheses include a unique zero-weight vacuum, positive integer weights for other symbols, and at most one weight-one symbol. Binary occupancy satisfies these hypotheses with weights 0 and 1. Its theorem also covers exact finite target reachability, but the present upper bound only needs finite halt patterns.

Corollary 27.1 (Sharp binary finite-input threshold). *In the class of fixed one-dimensional deterministic finite-radius binary conservative cellular automata with quiescent zero vacuum, five is the least particle count permitting a computable finite-input reduction of arbitrary Turing-machine halting to a designated finite anchored halt-pattern occurrence.*

Proof. The fixed universal construction above supplies the mass-five upper bound. If such a reduction used at most four particles, the decision procedure of Report 12 could decide the resulting finite-pattern instance and hence the original TM halting problem. This contradicts undecidability. The encoder is required to be computable, so it cannot hide the undecidable answer in an uncomputable choice of finite input. $\square$

[write] Corollary 16.1 of Report 15 (first half of this Part) implies this corollary: a globally reversible, number-conserving binary automaton is a binary conservative automaton, so its five-particle construction is also an upper bound here, and the lower bound is the same Report 12. The two upper constructions differ: Report 14 compiles every deterministic ADD/SUB/NOP program (no injectivity needed), and gives a literal universal instance, a radius linear in D and a persistent halt word; Report 15 needs partial-injective guarded sources and gives global reversibility and exact finite targets.

The all-rule mass-four classification is a separate mathematical dependency; this article does not claim to re-prove it. Its exact authored PDF, LaTeX source and required vector figure are included unchanged in the companion directory, with hashes. Its original executable replay remains in its separate release.

[write] Report 12 is printed as source 17 (Part VI) of `signal-machine-collision-certificates`; its byte-identical companion files delivered with Report 14 are not shipped here.

The sharp statement excludes infinite initial backgrounds, active zero-weight states, multiple unit labels, signed weights, spatially nonuniform or time-dependent rules, and other computational or observational models.

This matching of two proved bounds is not a claim of literature priority. Alhazov and Imai’s 2016 publisher abstract reports a five-particle universal numerical-state NCCA [5]; the full paper was unavailable in the underlying review. Kong’s 2021 binary thesis and Imai’s 2021 project report supply lower-particle precursors [6, 7]. A limited targeted search finding no reusable binary table or later resolution does not show that none exists. The present compiler is justified by its construction, and the lower bound by its companion proof, rather than by a present-day literature-status inference.

28 A source horizon polynomial compiler

Fix a source program and an instruction horizon $h \geq 1$ as part of the *compiler syntax*. This section gives a separate ordinary polynomial for each such h . For the universal instance, one unit of h means one literal ADD or SUB row of the 8,408-row two-counter program. It does not mean a TM transition, a virtual three-counter instruction, or a CA tick. These clocks are related by the explicit, input-dependent costs already proved.

[write] **Comparison.** For the same 8,408-row source, quadratic-orthant-certificates proves a stronger certificate, Theorem `qoc:um:thm:poly`: degree two, $29,904T$ witnesses and $4T + 1$ squares, exact over the naturals and over the nonnegative reals, where T is the instruction horizon. The polynomial of this section has degree at most four, $10,750h$ witnesses and $2,344h + 1$ squares, and needs a paid norm row for real witnesses (Section 29). Report 14 did not know that theorem. What this section adds is the optional row for the automaton’s clock, which ties the source horizon to the cellular-automaton time.

Expand each ADD and NOP row into one branch, and each SUB row into a positive decrement branch and a zero branch. Let B be the resulting number of branches. Thus $B = p + n + z$. Index branches by $j = 1, \dots, B$; let s_j, t_j be their integer source and target control codes under the injective ordering $Q \rightarrow \{0, \dots, m-1\}$. Let $\delta_{jk} \in \{-1, 0, 1\}$ be the branch's change to counter k . A zero or NOP branch has both changes zero. There are exactly z zero branches, and no branch starts at halt.

The initial control and natural counters q_0, a_0, b_0 are external inputs or fixed constants. The actual emitter in this release specializes them to supplied constants. The theorem-level formula also permits atomic input coordinates, but the current emitter does not export symbolic input wires.

28.1 All residuals and the witness variables

For every source layer $t = 0, \dots, h-1$, introduce B natural selector variables e_{tj} and two natural postcounter variables a_{t+1}, b_{t+1} . Use the following residuals, with $c_{t0} = a_t, c_{t1} = b_t$:

$$U_t = \sum_{j=1}^B e_{tj} - 1, \quad (28.1)$$

$$V_0 = \sum_j s_j e_{0j} - q_0, \quad (28.2)$$

$$V_t = \sum_j s_j e_{tj} - \sum_j t_j e_{t-1,j}, \quad 1 \leq t < h, \quad (28.3)$$

$$W_{tk} = c_{t+1,k} - c_{tk} - \sum_j \delta_{jk} e_{tj}, \quad k = 0, 1, \quad (28.4)$$

$$Z_{tj} = e_{tj} c_{tk(j)}, \quad j \text{ a zero branch}, \quad (28.5)$$

$$H = \sum_j t_j e_{h-1,j} - q_*. \quad (28.6)$$

Here $k(j)$ is the counter tested by zero branch j . Form the ordinary integer polynomial

$$P_h = \sum_{t=0}^{h-1} \left(U_t^2 + V_t^2 + W_{t0}^2 + W_{t1}^2 + \sum_{j \text{ zero}} Z_{tj}^2 \right) + H^2. \quad (28.7)$$

Every residual is explicitly emitted as a sparse integer polynomial. A monomial is a list of variable indices with repetitions denoting powers; an empty list denotes the constant monomial. No hidden integrality predicate, exponentiation atom, branch oracle or implicit guard is part of the polynomial. The baseline integrality condition is the stated witness domain $\mathbb{N}^N$.

Theorem 28.1 (Exact first-halt fiber). *For fixed natural external inputs and fixed syntactic $h \geq 1$, $P_h = 0$ has exactly one complete natural witness if the deterministic source first halts after exactly h literal instructions, and has no natural witness otherwise.*

Proof. A sum of real squares is zero precisely when every residual is zero. By nonnegativity and (28.1), each layer's natural selectors sum to one and therefore exactly one of them is one. The state equations and injective control codes force that selected branch's source to be the current control, starting at q_0 . The counter equations force its specified updates.

A selected zero branch has $e_{tj} = 1$, so (28.5) forces its tested old counter to be zero. A selected positive decrement at old value zero would force its postcounter to be -1 by (28.4), which is excluded by the natural witness domain. Thus no positive-test slack variable is needed. ADD and NOP branches have no additional guard. Since the source is deterministic, exactly one branch is

permitted at each nonhalting state and counter pair. Induction therefore fixes all selectors and postcounter values uniquely.

The terminal residual requires halt at layer h . It cannot have been reached earlier because halt has no outgoing branch, whereas each subsequent layer would require a branch with that source code. Hence a solution is precisely a first-halt trace of length h . Conversely, any such trace supplies its one-hot selectors and natural postcounters, making every residual vanish. Every coordinate of that trace witness is forced, proving uniqueness of the complete fiber, not merely uniqueness after projection to selected coordinates. $\square$

28.2 Exact dimensions and algebraic costs

The unsimplified core ledger is

$$\begin{aligned} \text{witness variables} &= h(B + 2), \\ \text{squared residual slots} &= h(4 + z) + 1, \\ \deg P_h &\leq 4. \end{aligned} \tag{28.8}$$

The degree-four bound comes only from squaring quadratic zero guards; specialization can lower actual degree. A residual that becomes identically zero remains a counted slot unless explicitly removed. Therefore slot counts and counts of nonzero residuals are reported separately.

With initial counters treated as atomic coordinates or constants, the written raw residual monomial occurrences before collecting like terms satisfy

$$S_{\text{raw}} \leq h(4B + 5) + 2. \tag{28.9}$$

Indeed the one-hot rows contribute $h(B + 1)$; state rows contribute at most $(B + 1) + 2B(h - 1)$; the two counter rows contribute at most $h(4 + B - z - n)$; the zero rows contribute at most hz ; and the terminal row at most $B + 1$. Their sum is $h(4B + 5 - n) + 2$, which gives the displayed bound. Zero coefficients, fixed input zeros and coincident terms can reduce it.

A sum-of-squares representation is an explicit representation of an ordinary polynomial, but expansion has a real cost. If the collected residual lengths are $\ell_1, \dots, \ell_s$, expanding all squares by ordered products takes at most

$$\sum_{r=1}^s \ell_r^2 \tag{28.10}$$

written term occurrences before global collection. The emitter reports the actual lengths' sum and this expansion bound. It never charges a square as though the fully expanded coefficients had no computational or storage cost. The arithmetic coefficients also have bit lengths; none of these counts is a bit-complexity claim.

28.3 Adding exact physical time

For branch j , let $\tau_j(a_t, b_t)$ equal one for a zero or NOP branch and $2c_{t,k(j)} + \kappa_j$ for an arithmetic branch, using (24.1). Add a natural witness coordinate Θ and one residual

$$T = \Theta - \sum_{t=0}^{h-1} \sum_{j=1}^B e_{tj} \tau_j(a_t, b_t). \tag{28.11}$$

Use $P_h + T^2$. Every summand selected by a valid trace is its proved CA duration, so Θ is exactly the first physical halt time. It is uniquely forced. The cost is one witness variable, one squared residual and degree still at most four. Alternatively Θ can be a specified input coordinate; then

there is no added witness variable, and the fiber is nonempty precisely when that time agrees with the first-halt clock.

If the source has no NOPs, the additional raw residual occurrences are at most $h(2B - z) + 1$: one selector term for each zero branch and two terms for each arithmetic branch, plus Θ . With NOPs the sharper expression is $h(2B - z - n) + 1$. Constant initial counters may combine the affine terms in the first layer and lower the actual count.

At $h = 0$, the witness-free core is simply $(q_0 - q_*)^2$; first physical halt time, when defined, is zero. The actual emitter's optional unsimplified clock version retains one Θ witness and two residual slots, $(q_0 - q_*)^2 + \Theta^2$. Substituting the known value $\Theta = 0$ removes that witness. Neither convention is silently charged as the other.

28.4 Numerical height and space bounds

Put $M = \max(a_0, b_0)$ and $\kappa_{\max} = \max(1, \max_q \text{ moving } \kappa_q)$, taking the inner maximum absent when there are no moving rows. In any valid h -instruction trace,

$$0 \leq a_t, b_t \leq M + h, \quad (28.12)$$

$$\Theta \leq h(\kappa_{\max} + 2M) + h(h - 1), \quad (28.13)$$

$$\text{diam} \left(\bigcup_{0 \leq u \leq \Theta} \text{supp } F^u(\text{Enc}(q_0, a_0, b_0)) \right) \leq 2Z + a_0 + b_0 + h. \quad (28.14)$$

For (28.12), a counter increases by at most one per source instruction. At layer t its old value is at most $M + t$, giving duration at most $\kappa_{\max} + 2(M + t)$ even for unit-time branches. Summation proves (28.13). For (28.14), every packet phase lies between the outer markers. The maximal leftward extension is at most the initial left extent plus the total number of ADDs to the left counter; similarly for the right. Across the *whole* trace, those two ADD counts sum to at most h . This establishes the stronger union-of-supports bound as well as a bound on each individual microstate's diameter.

Selectors have height at most one, so without the clock every witness coordinate is at most $\max(1, M + h)$; with the clock also take the maximum with (28.13). For the pinned source, independent exact enumeration gives

$$\min_q \kappa_q = 512,938, \quad \max_q \kappa_q = 529,754.$$

Thus its time bound uses 529,754 and its diameter bound starts at $807,240 + a_0 + b_0 + h$. These are numeric counter bounds, not polynomial bounds in the bit length of the original TM input. In particular the universal loader's $a_0 = 2^L 3^R$ remains explicit.

29 The paid nonnegative real variant

The natural-domain theorem cannot simply be read as a theorem over real witnesses. Nonnegative real numbers summing to one can have fractional selectors, and weighted source codes can then imitate an unrelated control. Pay for that domain change by adding one residual at each layer:

$$N_t = \sum_{j=1}^B e_{tj}^2 - 1. \quad (29.1)$$

The real-orthant polynomial is $P_h + \sum_t N_t^2$, together with T^2 when physical time is requested. All witness coordinates are now required to lie in $\mathbb{R}_{\geq 0}$, while the external initial counters remain natural.

Theorem 29.1 (Paid real-orthant fiber). *For fixed natural inputs and fixed source horizon, the paid polynomial has the same empty or singleton complete witness fiber over the nonnegative real orthant as the baseline polynomial has over naturals. It has the same core witness count, $h(5 + z) + 1$ core squared residual slots, and degree at most four.*

Proof. The one-hot sum and nonnegativity put every selector in $[0, 1]$. The additional norm identity gives

$$\sum_j e_{tj}(1 - e_{tj}) = 1 - 1 = 0.$$

Each summand is nonnegative, hence each selector is zero or one. Exactly one is one. Starting from natural inputs, the forced integer counter increments and decrements then make every postcounter an integer by induction; their nonnegativity excludes negative values. The proof of Theorem 28.1 now applies verbatim, and the optional exact clock is an integer forced by that trace. The norm residuals are quadratic, so their squares preserve the degree-four bound. There is one new square and $B + 1$ raw monomial occurrences per layer, with no new variable. $\square$

The additional rows are essential. Consider controls a, b, c with source rows $a : \text{ADD}(0, \text{HALT})$, $b : \text{ADD}(0, b)$ and $c : \text{ADD}(0, \text{HALT})$, ordered with codes 0, 1, 2 and halt 3. Start at b with $(a_0, b_0) = (0, 0)$ and $h = 1$. The true source never halts. Nevertheless the unpaid real relaxation admits $e_a = e_c = 1/2$, $e_b = 0$ and postcounters $(1, 0)$: the weighted source code is 1, the weighted target code is 3, and all counter equations hold. The added norm residual equals $-1/2$, so the paid sum has value at least $1/4$ and excludes this false halt. The independent audit checks this example using exact rational arithmetic.

This is a theorem about the closed nonnegative orthant, not unrestricted real witnesses or a strictly positive interior. It also retains the natural external-input restriction: the norm rows force discrete branch choices, but do not themselves force an arbitrary supplied real initial counter to be an integer.

[write] The norm row (29.1) is the same device as Report 15’s row (33) in Section 12.

30 Literal and small fully emitted examples

For the universal source, $B = 6,068 + 2 \cdot 2,340 = 10,748$. The parameterized core costs are therefore

Version	Witness variables	Squared residual slots
Natural core	$10,750h$	$2,344h + 1$
Natural with clock	$10,750h + 1$	$2,344h + 2$
Paid real core	$10,750h$	$2,345h + 1$
Paid real with clock	$10,750h + 1$	$2,345h + 2$

Degree is at most four in all four versions. These are source-parametric families of finite polynomials, with h fixed at emission time. They do not promote h to an unbounded input of one fixed polynomial, establish an unbounded single-fold Diophantine representation theorem, or set a small universal-polynomial record.

The file `universal_h1_frontend.json` emits every residual for the literal source at $h = 1$, source entry, input $(1, 0)$ and with clock. It has 10,751 witness coordinates, 2,346 squared residual slots, 1,176 nonzero residuals, 52,575 collected residual monomial occurrences and expansion bound 503,278,041 from (28.10). Because the initial counters are constant and there is only one layer, actual degree is at most two in this specialization. The first literal instruction does not halt,

and the fixture replay verifies that this schema has no accepting natural witness. The file is an explicit universal-source *schema instance*, not an exhibited accepting universal computation.

30.1 A complete accepting witness and CA trace

The small source in the replay is

```
start :  ADD(0, test),
test  :  SUB(0, HALT, bad),
bad   :  ADD(1, bad).
```

Its sorted nonhalt controls are bad, start, test. Consequently

	bad	start	test	HALT
$d(H_q)$	1	2	3	4
$d(O_q)$	5	6	7	
$d(I_q)$	8	9	10	

and $D = 10, C = 11, E = 22, K = 45, Z = 180, R = 337$. Start with $(a_0, b_0) = (0, 3)$. The literal source trace is

$$(\text{start}, 0, 3) \longrightarrow (\text{test}, 1, 3) \longrightarrow (\text{HALT}, 0, 3).$$

The first excursion lasts 237 ticks and the second lasts 235, so the first halt observation occurs exactly at tick 472. The source-boundary occupied sets are

$$\begin{aligned} t = 0 : & \quad \{-180, 0, 11, 13, 183\}, \\ t = 237 : & \quad \{-181, 0, 11, 14, 183\}, \\ t = 472 : & \quad \{-180, 0, 11, 15, 183\}. \end{aligned}$$

The complete direct sparse CA replay checks all 473 microstates, all five occupied sites at each one, the successor encodings, the absence of the 22-cell halt word before tick 472, and the fixed final configuration. Figure 2 is plotted from this verified trace data.

Order branches as bad ADD, start ADD, test positive SUB, test zero SUB. The two one-hot selector vectors are $(0, 1, 0, 0)$ and $(0, 0, 1, 0)$, with postcounters $(1, 3)$ and $(0, 3)$. In the emitted variable order, the complete clocked natural witness is

$$(0, 1, 0, 0, 1, 3, 0, 0, 1, 0, 0, 3, 472).$$

The file `example_frontend.json` explicitly contains its 12 squared residual slots, 11 of them nonzero, in 13 variables. It has degree at most four, 48 collected residual monomial occurrences and expansion bound 288. The first-layer zero guard specializes to zero since the initial tested counter is zero, explaining the one empty residual without changing the slot ledger. The provided witness evaluates the polynomial to zero. An independent fresh generator reconstructs both this JSON and its witness byte-for-byte, then separately replays the CA to validate that the polynomial clock is the actual first halt time.

30.2 External input variants and their costs

The universal raw positive-input interface can be parameterized mathematically by an external natural x using $a_0 = x + 1, b_0 = 0$. This is an input coordinate, not a trace witness. It preserves the variable, degree and uniqueness statements. However substitution into first-layer zero guards can produce additional monomial occurrences, for example $e(x + 1) = ex + e$. The atomic-input bound (28.9) must therefore be recomputed for that emitted substitution; it is not




Figure 2: Top: every occupied site at every tick of the accepting example. The two excursion clocks are 237 and 235. Bottom: two valid home configurations at control test, showing the fixed read site $-Z = -180$ for counter values zero and one. The home packet occupies 11 and 14 in both. The unchanged right endpoint at 183 is omitted from the bottom detail, and the broken axis compresses the remote empty interval. The orange circle marks a read location, not an additional particle or state.

inherited without charge. This optional symbolic substitution is not implemented by the present constant-input emitter.

Similarly the loading map $(L, R) \mapsto 2^L 3^R$ is outside the polynomial trace compiler. These exponentials are not silently treated as polynomial atoms inside (28.7). The certificate begins with the supplied natural physical counters and describes exactly h literal source instructions. The primary TM input, its stack encoding and prime loading remain separate computable inputs with potentially enormous costs.

31 Executable specification and independent checks

The released compiler `five_binary.py` is the executable counterpart of Sections 19–25; `frontend.py` emits the residuals in Sections 28–29. Their final SHA-256 digests are recorded in the release manifest and audit receipts:

```
five_binary.py
9022006a4669c7c637ac337a89ea792c55612c23a92c7bedda34af9c9707d90f

frontend.py
4d5b3ce9facb439d60a152890b3910c021b242996fbecbeefef7e4c0ba793ca5

source-replay/source/literal2.json
85e16b44828f2f3d4ad6d0805dcc9e9922893a6d286874f2018d6a33af864b00
```

The final API hardening changes neither the valid mathematical rule nor the small and universal polynomial fixture bytes. Source rows are copied into immutable snapshots; machine, CA and frontend lookup data cannot be changed after construction. Mutating the caller’s original dictionary or lists does not mutate the compiled rule.

Public source counters and horizons require exact nonnegative Python integers, including on halt fast paths; Boolean and floating-point substitutes are rejected. Occupied-site coordinates are exact integers, with negative integers allowed. Duplicates in a supplied coordinate iterable, Boolean coordinates and floating-point coordinates are rejected before converting to a set. Source instruction lengths, operations, counters and destinations are validated. These are host-language representation contracts. They must be distinguished from malformed *binary configurations*, whose distinct integer occupied sites are legitimate inputs to the CA and whose conservation was proved without an encoding promise.

The natural witness evaluator uses exact natural integers. The paid real version’s numerical sample evaluator accepts nonnegative integers and exact `Fraction` values, rejecting floats, NaN and infinities. The exported polynomial has the proved real-orthant semantics over all nonnegative reals; this exact sample evaluator is not a claim to implement arbitrary real arithmetic.

The evidence has four logically separate layers:

1. The mathematical component proof establishes arbitrary-configuration conservation, finite locality and all-natural-input simulation.
2. The source proof establishes the all-input TM-to-three-counter and three-to-two-counter simulation; the primary-source review establishes the finite-input universal machine and exact table interpretation.
3. Independently written finite tests challenge boundary geometry, local central outputs, conservation, exact clocks, observer exclusivity and witness fibers. They provide falsifiable regression evidence, not a replacement for unbounded proofs.
4. Fresh extracted replay verifies portable operation and byte identity of code, source, examples and report fixtures against the manifests.

The independent compiler audit covers 450 complete source macros, 52,885 analytical microstates, 238,835 local-center comparisons, 9,907 finite conservation cases and 9,520 translation comparisons. It also includes 1,190 pair/triple shapes with sensor variants, 4,096 exhaustive dense words, 51 long connected boundary phases, 1,000 adversarial multiple-component cases and 288 periodic words. Separate literal-source boundary replay covers every one of the 8,408 controls, 16,816 zero/one counter cases and 205,004 transitions at dispatch, separation, target contact, reversal and home commit.

The frontend audit checks 2,016 fixed-input instances, 27,504 one-hot branch sequences, 888 rational-coordinate mutations and 4,998 rational simplex points. It checks the unpaid fractional counterexample and the paid correction. The normal and optimized API suites each reject 904 invalid calls and survive 24 mutation/deletion attempts, including caller snapshot isolation. The full independent mathematical regression suites run against both ordinary and actual Python -O compiler imports. In the optimized runs, the audit entry points are compiled with assertions explicitly retained; an optimized interpreter therefore cannot make the audit pass by discarding its checks.

The producer also records 252 source macro cases, 186,386 CA microsteps, 428 multi-instruction source steps, 1,200 malformed finite configurations, 24,769 local-site checks and 20 local boundary cases. These ledgers are reproducible counts of executed checks. None is a machine-checked theorem about all source programs or all binary neighborhoods.

31.1 Portable replay and artifact identity

The top-level `replay.py` requires only Python 3.10 or later and its standard library. It verifies all manifest hashes, runs the source verifier, producer checks, independent normal and optimized suites, API checks, and report-fixture reconstruction, and then verifies unchanged manifest bytes again. It also tests the mathematical values printed in this report against the exact fixture JSON. The release verification receipt records the fresh archive extraction and replay result. No network or package installation is needed for replay.

The PDF was built from the accompanying LaTeX and all pages were rendered and visually inspected. Rebuilding the main document requires a conventional LaTeX installation with the listed packages and the supplied vector figure. Re-rendering the optional figure requires matplotlib; figure regeneration is not a dependency of the standard-library verification. The vector graphic, its source and the exact 473-state trace data are all included and hashed. The companion Report 12 files are byte-preserved and independently hashed, not regenerated as part of the new replay.

[write] Neither the companion Report 12 files nor the delivered PDF are shipped with this report; the replay files are shipped under the prefix `14-five-binary-`, and three heavy generated inputs of the replay (`universal_h1_frontend.json` and the two `literal2` files) are rebuilt as described in the README.

32 Further research directions

The following are useful directions suggested by the construction. They are not assertions that no prior work addresses them.

1. **Reduce the geometric constants.** The buffer inequalities are intentionally generous. A smaller K or a tighter locality proof could reduce radius and observer length. A change must preserve malformed-input conservation, not only valid trajectories, and must re-check every contact phase.

2. **Compress control without extra particles.** The present packet gap alphabet has $m + 2p$ codes. Determine which control and phase distinctions can be merged without ambiguous travel direction, incorrect zero branches or a premature halt code.
3. **Obtain alternative literal universal sources.** Fewer ADD/SUB controls directly reduce D , R , the indexed schema count and polynomial branch count. Any replacement needs its own finite-input interface, all-input proof and complete pinned literal table; row count is not a universality test. [write] Report 16 (Part II) supplies another literal universal source with a complete pinned table and all-input proof, reversible but much larger (122,622 controls, 141,561 rows), for Report 15's reversible compiler; it does not reduce the constants of this construction.
4. **Formalize the full-shift extension.** A proof-assistant development could link component parsing, disjoint bounded-write hulls and the local Boolean evaluator, then verify the all-input simulation invariant and exact clock. This would close the current gap between mathematical audit and formal proof.
5. **Account for bit complexity end to end.** The current report charges numeric counter height, physical time, spatial extent and algebraic occurrences. A full bit-cost analysis should include stack loading, prime encoding, arbitrary-precision arithmetic, source expansion and polynomial serialization.
6. **Separate finite-horizon from unbounded arithmetic.** The explicit unique-witness family does not combine its unbounded horizon into one fixed polynomial. Investigate what additional machinery would achieve such a theorem, without treating a family of ever-larger tableaux as a fixed-arity result.
7. **Study other particle models and observations.** Several unit labels, active zero-weight backgrounds, reversible restrictions and intrinsic universality alter the assumptions behind the lower or upper bounds. Their thresholds require separate arguments.
8. **Complete the primary literature comparison.** Obtain missing full papers and later primary work before making any novelty, priority or optimal-resource claim. The construction's internal correctness does not answer that historical question.

Part II

A literal universal reversible source (Report 16)

A literal universal reversible two counter source

Complete expanded transitions and all input proofs for a binary five particle compiler

Report 16 Mathematical construction and reproducible audit

3 October 2026

Abstract of Report 16, as delivered. This report supplies the fully expanded universal source left existential in Report 15. One pinned machine has 122,622 controls and 141,561 literal zero-test, positive-test, increment, decrement and identity branches on exactly two natural counters. Its transition function is a partial injection on all natural configurations, its start control has no incoming branch, and its halt control has no outgoing branch. We give the finite-tape loader, complete unbounded-input invariants and termination ranks for every simulation layer, and exact state-dependent clocks. The construction starts from the published Neary–Woods 15-state two-symbol Turing machine and uses directly proved and independently reconstructed reversible history and prime-arithmetic graphs. Its clean input family has no premature blocking: nonhalting is equivalent to infinite continuation. Consequently the fixed binary five-particle compiler instance also has undecidable positive return and periodicity on this family. All source rows, reconstruction checks and source-level ledgers are bundled. The much larger cellular-automaton factor list, local truth table and universal-instance quartic polynomial are not materialized. Even empty-tape initialization has a predicted two-counter clock near $8 \cdot 10^{13}$ steps; only small literal macros and the fresh entry have been executed. A shared-offset primitive certificate reduces the fixed-horizon core to $141\,565K$ witnesses and $23\,435K + 1$ squares. Exact certificate counts, reproducibility boundaries, and useful next questions are included. This is an ordinary mathematical proof with executable audits, not a proof-assistant formalization, efficient simulation, novelty claim or size record.

[write] Delivered as `Literal_Universal_Reversible_Source_Package.zip` (batch 82, archive 13; arrival `db37d18c8`), dated 3 October 2026; files placed by `bca6383e9` with the prefix given in Appendix A. Report 16 materializes the universal source that Report 15 (Part I) left existential. Its main file `report16.tex` and its ten section files are inlined.

33 The machine and the exact claim

We use $\mathbb{N} = \{0, 1, 2, \dots\}$. An instantaneous description (ID) of a k -counter machine has a control $q \in Q$ and a vector $(c_0, \dots, c_{k-1}) \in \mathbb{N}^k$. A primitive (q, i, x, t) uses a counter index i , source and destination controls q, t , and one of five symbols:

Symbol	Enabling condition	Counter change	Cost
Z	$c_i = 0$	none	1
P	$c_i > 0$	none	1
$+$	true	$c_i \leftarrow c_i + 1$	1
$-$	$c_i > 0$	$c_i \leftarrow c_i - 1$	1
0	true	none	1

Successful tests are separate steps. A combined conditional decrement is used only in an earlier irreversible layer and is explicitly split before the reversible construction. No unbounded register is stored in a control name.

Theorem 33.1 (Pinned literal universal source). *Let M be precisely `source/source.json` in this release. It has 122,622 controls and 141,561 branches, two natural counters, and class cutoff $J = 0$. Its one-step relation is an injective partial function on $Q \times \mathbb{N}^2$. No branch enters **START**, and no branch leaves **HALT**.*

Let U be the Neary–Woods machine in Table 16 of [1]. For every $L, R, T \in \mathbb{N}$ and every $C \geq 1$ with $\gcd(C, 2310) = 1$, the run of M from

$$(\text{START}, C2^L 3^R 5^T, 0) \tag{1}$$

*reaches **HALT** if and only if U , started in state $A = u_1$, scanning the blank symbol $0 = c$, with finite half tapes represented by L, R , reaches its undefined $J1 = (u_{10}, b)$ transition. Every such source run either reaches **HALT** in finite time or continues forever. This input family includes an effective encoding of every program and finite input in the published universality theorem. Thus this one fixed literal machine is computation universal.*

The theorem’s proof is the complete layer-by-layer argument below. The table’s SHA-256 is

38c586706fa1069442d3e5103151adb156448
b7d7c783f5c46d9d507b87fe73a

The line break is typographical; the digest has 64 hexadecimal characters. The release manifest pins every other artifact independently.

The unrestricted all-natural statement is *partial injection*; the universality and no-premature-blocking statements have the clean-input promise (1). An arbitrary second physical counter is not permitted. Nor may arbitrary initial powers of 7 and 11 be silently added: these encode history and work counters in the construction. The advertised loader initializes both to zero. The history invariant actually permits an arbitrary initial history exponent, but does not permit arbitrary nonzero work. We retain the narrower clean API.

33.1 What is literal and what remains a finite specification

Every operation of M is a row in the saved JSON. The Python generator, macro certificates, prime list, proofs and checkers are *not* primitives available to M . They explain or verify the finite table. There are no row callbacks, arithmetic oracles, unexpanded macro invocations, symbolic infinite lists or hidden work counters in the two-counter source.

By contrast, the downstream binary CA is given by the already audited finite *indexed construction* of Report 15 applied to this literal source. Its enormous factor array and truth table have not been allocated. The fixed source horizon quartic is likewise specified by a counted finite formula; no giant universal-instance polynomial is included. Distinguishing these materialization levels prevents the new source artifact from implying work that was never performed.

34 The finite tape and the irreversible front end

[write] **A second presentation.** The 528-instruction three-counter front end of this section, with its file `virtual3.json`, is byte for byte the program of quadratic-orthant-certificates Part V (shipped there as `data/16-universal-membrane-virtual3.json`); its correctness is that




Figure 3: All counter-machine stages are saved as finite literal tables. The last stage, rather than an existence citation, is the input to Report 15.

report’s Lemma `qoc:um:lem:virtual`. It is kept here because the later layers cite its clocks. The prime encoding of Section 36 is a reversible variant of that report’s Lemma `qoc:um:lem:prime`, not the same lemma.

34.1 Primary table and exact finite loader

The 30 cells of Table 16 in [1] were compared with the saved `tm_table.json` against rendered primary PDF page 17 (printed page 121). Printed page 112 was also visually checked for the input conventions. The machine table, using $0 = c$, $1 = b$, and states $A, \dots, O$ for $u_1, \dots, u_{15}$, is reproduced here. An entry writes its first symbol, moves as specified, and enters its last state.

Scanned	A	B	C	D	E
0	0RB	1RC	0LG	0LF	1RA
1	1RA	1RA	0LE	1LE	1LD
Scanned	F	G	H	I	J
0	1LD	0LH	1LI	0RA	1LK
1	1LD	1LG	1LG	1LJ	undefined
Scanned	K	L	M	N	O
0	0RL	0RM	0LB	0LC	0RN
1	1RN	1RL	1RL	0RO	1RN

The undefined cell is $J1$, not $J0$. A known prose/display discrepancy later in the article is not a reason to alter the pinned transition table. The old bibliographic comment in the byte-preserved compact machine file gives inaccurate page numbers; the bibliography of this report uses the corrected 105–126 range without editing that dependency.

[write] **Correction (3 October 2026)**. Here too the “old bibliographic comment” gives the publisher’s pages: Neary–Woods is *Fundamenta Informaticae* 91(1) (2009), 123–144, DOI 10.3233/FI-2009-0036; 105–126 is the author-hosted PDF’s pagination, and the DOI 10.3233/

FI-2009-0008 that Report 16 printed belongs to another paper (Appendix A). The table above agrees entry by entry with Table 5 and with quadratic-orthant-certificates' Table `qoc:wf:tab:tm`, where the (u_{10}, b) discrepancy is recorded.

Definition 3.1, equations (3)–(4), Tables 1–3 and Section 3.5 of the primary paper specify the finite program/data encoding. Its separator G for this machine is bc , and the initial head is on its rightmost symbol $c = 0$. Thus fixing the initially scanned symbol to zero retains the universal input family. Both tails are blank c , rather than nonblank periodic backgrounds. The reduction from an arbitrary program/input to that finite universal tape is the published universality dependency, not a claim that the loader below reimplements the full program compiler.

Place the head at coordinate zero. For finite nearest-head-first words $\ell = \ell_0 \cdots \ell_{a-1}$ and $r = r_0 \cdots r_{b-1}$, put symbol ℓ_i at $-i - 1$ and r_i at $i + 1$. Unmentioned cells and the scanned cell are zero. Define

$$L = \sum_{i < a} \ell_i 2^i, \quad R = \sum_{i < b} r_i 2^i, \quad \mathcal{L}(\ell, r) = (\text{START}, 2^L 3^R, 0). \quad (2)$$

`loader.py:from_tape` implements exactly (2). Its checks reject nonbinary characters and non-string tape arguments. The counter interface `from_counters` supplies optional T, C from (1), rejecting Boolean pseudo-integers, negative values, zero cofactors and cofactors not coprime to 2310. These APIs return exact integers. The command-line interface prints hexadecimal to avoid interpreter decimal-digit limits.

Every input gives finite values by finite sums and powers. This is a computability statement, not a feasible-memory promise. For example, `left=101,right=01` means $L = 5, R = 2$ and physical input $A = 288, B = 0$. The loader neither simulates U nor solves any halting question.

[write] Added 3 October 2026, batch 82 (reciprocal note): Report 32, printed as Part V of `group-theoretic-substrates`, cites this Report's finite $U_{15,2}$ input convention, its loader and machine table, in its delivered dependency audit, for the input hardness of its fixed matrix semigroup only (Theorem `gts:um:thm:semigroup` there).

34.2 Three counters with the scanned symbol in control

The pinned three-counter program has 528 active instructions: 295 increments and 233 combined conditional decrements. It uses data counters L, R and scratch T . A combined instruction $\text{SUB}(i, q, z)$ decrements a positive c_i and goes to q , or leaves zero unchanged and goes to z . An increment is unconditional. Consequently every nonhalt instruction is enabled on every natural vector.

At a TM cut, the state and scanned bit are in finite control, the two half tapes use low-order bits nearest the head, and $T = 0$. The initial prologue is

$$\text{init_clear_T} : \quad \text{SUB}(T, \text{init_clear_T}, \text{tm_A0_pop0}). \quad (3)$$

It takes exactly $T + 1$ combined instructions and clears arbitrary initial scratch. The later reversible history retains enough information to invert this clearing; no irreversible erasure is asserted for the final machine.

For a defined TM instruction $(q, s) \mapsto (w, D, q')$, let X be the half tape in the direction D and Y the other half tape. Write $X = 2Q + \rho$, $\rho \in \{0, 1\}$. The following finite graph is instantiated separately for each instruction and parity branch. Labels denote controls, not procedures called by the final machine:

$$\begin{aligned} \text{pop0} &: \text{SUB}(X, \text{pop1}, \text{back0}), \\ \text{pop1} &: \text{SUB}(X, \text{pair}, \text{back1}), \\ \text{pair} &: \text{ADD}(T, \text{pop0}). \end{aligned}$$

After k full loops, $X = 2(Q - k) + \rho$ and $T = k$. For $k < Q$ both subtractions are enabled. The rank $Q - k$ decreases. At $k = Q$, one failed test if $\rho = 0$, or one success and a failure if $\rho = 1$, determines the bit ρ and leaves $X = 0, T = Q$.

For the selected parity, a transfer loop alternates a successful $\text{SUB}(T)$ with $\text{ADD}(X)$, then takes the zero exit. After j transfers $X = j, T = Q - j$; rank T proves termination at $X = Q, T = 0$. The whole pop stage takes

$$3Q + 1 + \rho + 2Q + 1 = 5Q + \rho + 2. \quad (4)$$

The push loop consumes one Y and adds two units to T per iteration. After k iterations $Y = Y_{\text{old}} - k, T = 2k$; rank Y decreases. At exit, $Y = 0, T = 2Y_{\text{old}}$. A second transfer loop decreases T one at a time and increases Y , preserving $Y + T = 2Y_{\text{old}}$. This terminates at $Y = 2Y_{\text{old}}, T = 0$. If $w = 1$, one final increment writes the new low bit; if $w = 0$, the control edge is linked directly to the next cut. The next state is (q', ρ) . The push cost is $7Y_{\text{old}} + 2 + w$. Hence one TM step costs exactly

$$5Q + \rho + 7Y_{\text{old}} + w + 4 \quad (5)$$

combined instructions and realizes $(X, Y, T) \mapsto (Q, 2Y_{\text{old}} + w, 0)$, exactly the TM update.

All 29 defined TM cells use this graph. The undefined $J1$ continuation is **HALT** and has no instruction. The loops have no internal halt and terminate for every natural half tape by the displayed ranks. A cut requires both its control label and $T = 0$: the pop loop can revisit its entry label with $T > 0$. Checking labels alone would give a false simulation boundary. The saved-row checker covers all 528 instructions with 437 affine paths, using arbitrary natural parameters rather than a finite sample grid.

34.3 Separate tests and a genuinely fresh start

Replace every combined $\text{SUB}(i, q, z)$ by a zero test to z , and a positive test to a fresh private control followed by a decrement to q . The zero case takes one primitive step; the positive case takes two. An original increment still takes one. If a virtual trace uses v combined instructions with v_+ successful conditional decrements, its split cost is $v + v_+$.

The old initial control in (3) has a self-loop and cannot be declared predecessor-free. Insert **START** with one identity edge to that control *before* reversibilization. This changes no data and costs one step. The literal split table has 763 controls and 995 rows. Its start has no incoming row and halt has no outgoing row. No reversible edge is casually prepended at a later stage, where it could create an overlapping range.

34.4 Reducing indegree without changing data

For an original target with ordered incoming rows $e_1, \dots, e_k$, $k > 2$, introduce $k - 2$ private merge controls $u_1, \dots, u_{k-2}$. Redirect e_1, e_2 to u_1 , and e_j to u_{j-1} for $3 \leq j < k$, leaving e_k directed to the old target. Chain the new controls by identity rows toward that target. Every new merge has indegree two, and the old target receives only the chain's last edge and e_k . No operation is changed.

The extra cost after row e_j is

$$k - \max(2, j). \quad (6)$$

The number of remaining chain nodes is a strictly decreasing rank; all chains are finite and cannot block. The present table has 16 affected targets, adding 29 controls and 29 identity rows. The normalized source has 792 controls and 1,024 rows. Edges into halt are normalized too. The no-incoming start and no-outgoing halt conditions survive unchanged.

35 Reversible history with five counters

[write] Sections 35–37 make literal the existence argument of signal-machine-collision-certificates Part IV, §smc:s1:sec:safe (Morita’s history counter, a scratch counter, prime encoding with a scratch-zero invariant, and induction on legal progress): here the machine is fixed and its costs and termination ranks are exact. In substance this answers that report’s question “Print one safe universal source” (Part IV, §smc:s1:sec:questions, item 1), for the guarded syntax of Report 15 rather than source 13’s pulse interface.

The construction is inspired by Morita’s separate-test and separate-move machines [2], in particular the history and counter-reduction constructions. Primary text extraction was consulted at printed pages 308–309 and 314–316. The scan could not be visually verified: PDF screenshot requests failed and a live download endpoint returned HTML. We therefore make no visual-transcription claim for Morita’s diagrams or ambiguous OCR. Instead the actual local graphs below are explicitly generated, independently reconstructed from their preceding layer, and proved directly. No uncertain OCR symbol is a premise of the proof.

Add history H and work W to the normalized three-counter source. Its boundary invariant is

$$(\mathbf{d}, H, W) = (\mathbf{d}, h, 0), \quad h \in \mathbb{N}, \quad (7)$$

where $\mathbf{d} = (L, R, T)$. At a target with fewer than two incoming edges, or with a complementary Z/P pair on the same counter, leave the incoming rows unchanged. Those images are already disjoint. Such a row costs one step and leaves H, W unchanged.

Every other indegree-two target has a colliding pair. Give its rows branch numbers $b = 0, 1$. Each new path first executes its original source operation, then checks $W = 0$ and enters a history transfer loop. In particular, a negative update is blocked before history work when its original counter is zero. The source guard has not been weakened.

Lemma 35.1 (Exact history macro). *For every enabled colliding source row and every $h \geq 0$, its expanded path maps $(\mathbf{d}, h, 0)$ to $(\mathbf{d}', 2h + b, 0)$ at the old destination in exactly $10h + 4 + 2b$ primitive steps. It cannot stop or run forever internally.*

Proof. After the source operation and work-zero guard, the transfer loop has

$$H = h - j, \quad W = j, \quad 0 \leq j \leq h. \quad (8)$$

One iteration tests $H > 0$, decrements H , increments W , and tests $W > 0$. It costs four steps. Rank H decreases, and every decrement has its required positive test. At the final $H = 0$ test, the pair is $(0, h)$.

Branch zero enters the doubling loop directly. Branch one first increments H and tests $H > 0$, taking two steps. At the doubling-loop test after j iterations, the invariant is

$$H = b + 2j, \quad W = h - j, \quad 0 \leq j \leq h. \quad (9)$$

One iteration tests $W > 0$, decrements W , increments H , tests $H > 0$, increments H again and tests $H > 0$ again. It costs six steps and decreases rank W . The final zero test on W reaches the original destination with $(2h + b, 0)$. All post-increment positive tests are enabled. The exact cost is

$$1 + 1 + 4h + 1 + 2b + 6h + 1 = 10h + 4 + 2b. \quad (10)$$

Both ranks and all guards remain valid when $h = 0$: the two branches cost four and six steps respectively. $\square$




Figure 4: The two history invariants. Branch one pays two additional steps between the loops. This diagram summarizes the fully expanded finite graph; it is not an extra machine operation.

35.1 Why history recording is globally reversible

The guards inside this graph do more than ensure intended execution. At a transfer-loop entry, its first arrival has $W = 0$, while loop returns have $W > 0$. At the joins in the doubling part, the relevant incoming paths have $H = 0$ and $H > 0$. These are disjoint image conditions on all natural inputs. Replacing the displayed tests by identity arrows would destroy reversibility.

More generally, the independent graph checker groups all rows by source control and all rows by destination control. Any group with more than one row must be exactly a complementary Z/P pair on the same counter. Each single primitive is an injective partial map; those pairs have disjoint domains when outgoing and disjoint images when incoming. Therefore the *entire* five-counter relation is deterministic and partial-injective, including on unclear work values or unreachable configurations. This is a finite structural proof of an all-natural property, not a test of sample histories.

There are 233 recording pairs, hence 466 replaced incoming rows. Each pair's explicit graph uses 21 rows in place of its two and adds 16 controls. Thus

$$1024 + 233(21 - 2) = 5451, \quad 792 + 233 \cdot 16 = 4520. \quad (11)$$

The actual row multiset is compared with an independently reconstructed graph; the count is not used as a substitute for that comparison. Private controls are fresh and no source destination is entered before the macro terminates. All normalization edges, including those into halt, were included in this history pass.

36 Prime encoding into exactly two counters

Use primes $(p_0, p_1, p_2, p_3, p_4) = (2, 3, 5, 7, 11)$. At a five-counter boundary, represent $\mathbf{a} = (a_0, \dots, a_4)$ by

$$A = C \prod_{i=0}^4 p_i^{a_i}, \quad B = 0, \quad C \geq 1, \quad \gcd(C, 2310) = 1. \quad (12)$$

Unique factorization says $a_i > 0$ exactly when p_i divides A ; it says $a_i = 0$ exactly when p_i does not divide A . The cofactor C is preserved. Although A can vanish inside a macro, it is positive at every boundary. At the initial boundary, $a_3 = a_4 = 0$, giving (1).

One identity source edge is copied as one physical identity edge. For all other source edges, a zero test first checks $B = 0$. The remainder below is always a bounded finite-control index from 0 to $p - 1$, not an uncounted natural register. The construction has 4,519 source macros, one for each nonterminal five-counter source control, and 2,330 destination-shared restorers. Complementary

source tests at one control share one division macro. All graphs are expanded in the saved two-counter table.

36.1 Common transfer for arithmetic moves

For a prime p and boundary value $A = N \geq 1, B = 0$, both the multiplication and division graphs first transfer A to B . At the loop test after j iterations,

$$A = N - j, \quad B = j, \quad 0 \leq j \leq N. \quad (13)$$

An iteration tests $A > 0$, decrements A , increments B , then tests $B > 0$. It costs four steps and decreases rank A . The final $A = 0$ test produces $(0, N)$. Including the initial $B = 0$ guard, this phase costs $4N + 2$. The post-increment positive guard is important to the all-natural reverse join, even though it is automatic on the intended forward path.

36.2 Multiplication and an encoded increment

At the multiplication-loop test after j iterations,

$$A = pj, \quad B = N - j, \quad 0 \leq j \leq N. \quad (14)$$

One iteration tests $B > 0$, decrements B , performs exactly p increments of A , then tests $A > 0$. It costs $p + 3$. Rank B decreases and the final $B = 0$ test reaches the five-counter target encoded as $(pN, 0)$. The full clock is

$$4N + 2 + (p + 3)N + 1 = (p + 7)N + 3. \quad (15)$$

Every decrement is enabled and every positive post-increment test is valid; there is no internal blocking on a promised input. Multiplying by p_i increments just the corresponding encoded exponent.

36.3 Division and an encoded decrement

For an enabled source decrement, $N = pQ$ with $Q \geq 1$. After the common transfer, the division-loop invariant after j groups is

$$A = j, \quad B = p(Q - j), \quad 0 \leq j \leq Q. \quad (16)$$

A group tests $B > 0$, makes exactly p decrements of B , increments A , and tests $A > 0$. It costs $p + 3$. Since B is a positive multiple of p at each continuing test, all p decrements are enabled. The rank $Q - j$ decreases. The final zero test returns $(Q, 0)$, with exact total cost

$$4N + (p + 3)Q + 3 = (5p + 3)Q + 3. \quad (17)$$

This decreases precisely the chosen exponent. On an unpromised $N = pQ + \rho$, $0 < \rho < p$, the graph makes Q full groups, consumes the remaining ρ units, and blocks at a forced decrement before a source boundary. Successful termination is not claimed on this disabled case. In particular it does not falsely reach halt. Global partial injection of these malformed runs is a separate structural property proved below.

36.4 Division with a bounded remainder for tests

A test source control has either one Z or P row, or a complementary pair on one counter. Let $N = pQ + \rho$, $0 \leq \rho < p$. At division control j after k complete groups, the invariant is

$$A = N - pk - j, \quad B = k, \quad 0 \leq j < p, \quad (18)$$

whenever the right side is nonnegative. Each consumed unit of A requires a positive test and decrement, taking two steps. Every full group closes with an increment of B and a positive test on B , adding two steps. The unconsumed A strictly decreases between at most two group-close steps. Thus no infinite run is possible in the division phase. Equivalently one may use a lexicographic rank consisting of unconsumed units and a finite pending phase index.

The first zero test on A occurs at remainder control ρ , with $A = 0, B = Q$. The clock through that zero test, including the entry guard, is

$$2N + 2Q + 2. \quad (19)$$

Remainder zero means the encoded counter was positive and selects its P destination. Remainders $1, \dots, p-1$ mean it was zero and select its Z destination. If the source has only one of these tests, the other exit is absent. A false singleton test blocks in this division graph and never enters a source boundary. For an enabled source test, the correct exit exists.

36.5 Destination sharing is essential for restoration

For the selected target, enter its *shared* restorer at remainder index ρ . First add ρ to A , using an increment and positive test for each unit. At remainder-zero control the pair is (ρ, Q) . After j whole restoration groups,

$$A = \rho + pj, \quad B = Q - j, \quad 0 \leq j \leq Q. \quad (20)$$

A group consists of a positive test and decrement on B , then p pairs of increment and positive test on A . Its cost is $2p + 2$ and rank B decreases. The last zero test returns $(N, 0)$ to the chosen five-counter destination. Restoration costs

$$2\rho + (2p + 2)Q + 1. \quad (21)$$

Combining (19) and (21) and using $N = pQ + \rho$ gives the exact test clock

$$4N + 4\lfloor N/p \rfloor + 3. \quad (22)$$

Tests preserve the entire exponent vector and cofactor.

It would be incorrect to build independent restorers per incoming test edge: two copies would finish with overlapping $B = 0$ images at a common source destination. Instead restoration is shared by destination. Global reversibility of the five-counter graph ensures that tests entering one such destination use the same source counter and have at most one zero and one positive row. The distinct remainder paths enter the shared graph through complementary zero/positive guards. The independent reconstruction checks those exact joins and every restorer's prime and target.

Lemma 36.1 (Prime macro correctness and exact costs). *Every enabled five-counter row on a boundary (12) has a unique finite physical path to its correct encoded destination. No other five-counter boundary is visited internally. With selected prime p and pre-step encoded integer N , its cost is*

<i>Source operation</i>	<i>Physical primitive steps</i>
<i>identity</i>	1
<i>increment</i>	$(p + 7)N + 3$
<i>enabled decrement</i>	$4N + (p + 3)(N/p) + 3$
<i>enabled zero or positive test</i>	$4N + 4\lfloor N/p \rfloor + 3$

Every cost is a finite positive integer. The cofactor remains unchanged.

Proof. The preceding invariants prove the stated postconditions and exact clocks. Their ranks prove termination, and the positivity arguments prove that no required decrement or test blocks. All macro-internal controls and restorer controls are fresh relative to the five-counter boundary controls. The final edge alone reaches the intended source destination. Unique factorization translates multiplication, division and remainder tests into exactly the chosen exponent operation. $\square$

36.6 Partial injection on every pair and strict serialization

The two-counter graph satisfies the same structural criterion as the five-counter graph: any multiple outgoing or incoming group consists of complementary Z/P tests on one and the same counter. Thus primitive domains and images are pairwise disjoint on every natural pair. Every single primitive is individually injective. This proves global partial injection without a well-formedness promise.

The strict serialized schema maps counter index zero to side -1 , index one to side $+1$, updates to $-1, 0, +1$, and guards as follows:

Primitive	Serialized domain	Exact image condition
$+$	true	selected counter > 0
$-$	selected counter > 0	true
Z	selected counter $= 0$	selected counter $= 0$
P	selected counter > 0	selected counter > 0
0	true	true

All other counters are unchanged. Hence domain and image predicates are constant on each of the four cells formed by $\{0, > 0\}^2$; $J = 0$ suffices. The indexed validator derives each image from its original update and guard, including nonnegativity of the predecessor. It does not trust a separately supplied image predicate. It checks exact field names and types, valid references, unique branch names, the guard syntax, natural outputs, source and target class disjointness, no incoming start and no outgoing halt. Its four-cell proof covers all natural pairs, rather than only four tested sample values, because each predicate is constant on those cells.

The final table has 33,436 increments, 32,630 decrements, 23,429 zero tests, 52,036 positive tests and 30 identities. Their sum is 141,561. Thus the compiler moving and zero-update row counts are

$$p = 66\,066, \quad a = 75\,495. \quad (23)$$

Certificates specify private-state names, but independent reconstruction checks their claimed source rows, targets, primes, kinds and coverage before comparing complete edge multisets. Certificates are not trusted assertions that an unchecked macro worked.

37 Composition on all promised inputs and exact clocks

37.1 Finite prefixes and the no blocking theorem

At an intended boundary, each simulation layer agrees with the preceding layer on data and control. Section 34 proves the exact TM transition for arbitrary natural half tapes and clears arbitrary initial scratch. Splitting preserves each branch with one or two primitive steps. Normalization follows it by its finite identity chain. The history layer preserves the three data counters and updates only (H, W) as in Lemma 35.1. The prime layer preserves the representation (12) and realizes each five-counter primitive by Lemma 36.1.

Initialize (L, R, T, H, W) to $(L, R, T, 0, 0)$. For every finite valid source prefix, induction on its transitions shows that the literal simulation reaches the corresponding encoded boundary after the sum of the finite positive macro costs, with the stated invariants. Fresh internal controls prevent an unintended layer boundary from being mistaken for the final state. In the three-counter front end the additional $T = 0$ cut condition is used wherever the loop revisits its entry label. All edges into halt are included in every expansion; no premature internal halt is available.

Proposition 37.1 (Clean runs cannot block before halt). *On every input (1), either M first reaches HALT in finite time or its literal primitive run is infinite. The second case occurs exactly when U does not halt on the associated finite tape.*

Proof. The scratch prologue terminates by decreasing T . Every nonhalt combined three-counter instruction is enabled on every natural vector. Every defined TM step expands into a finite terminating front-end macro. Every enabled split/normalized edge expands to a finite enabled history macro on $W = 0$; every resulting five-counter edge expands to a finite enabled prime macro on (12). Thus a valid boundary cannot lead to nonfinal blocking or an infinite internal stutter. A finite TM halting path reaches literal halt by composition. If U never halts, it has infinitely many defined steps, and therefore infinitely many literal steps, since every expanded duration is positive. Conversely a literal halt must occur at the correctly simulated halt boundary. This proves both alternatives and completes Theorem 33.1. $\square$

The proposition is stronger than halting equivalence alone. An arbitrary guarded source could have a finite nonhalt stuck run, which still reflects periodically under Report 15. Such an input is ruled out here only on the clean family (1). It is not ruled out for every malformed pair in $\mathbb{N}^2$.

37.2 The complete clock composition

There is no claimed uniform constant slowdown. Use the actual intermediate values to compute the following exact clock layers:

1. The initial scratch prologue takes $T + 1$ combined instructions; each TM step takes (5) with its own $Q, \rho, Y_{\text{old}}, w$.
2. Add one fresh-entry step. Charge one for each increment or zero conditional-decrement branch, and two for each positive conditional-decrement branch.
3. For each primitive edge e_j entering a normalized indegree- k target, append the $k - \max(2, j)$ identity steps in (6). Unchanged targets append zero.
4. Let h_s be the pre-step history of each normalized edge. An unchanged edge has cost one and $h_{s+1} = h_s$. A recording edge of branch b_s has cost $10h_s + 4 + 2b_s$ and $h_{s+1} = 2h_s + b_s$.

5. For *every* resulting five-counter primitive, including history and work operations, form its actual encoded pre-step integer N_t from (12), choose its counter's prime, and sum the costs in Lemma 36.1. This gives the exact two-counter primitive time.
6. For every literal two-counter row of the resulting run, use the Report 15 CA clock in Section 38. This last sum is the exact physical CA observation time. One CA step means its entire fixed factor composition, not one individual local involution.

Each item is a finite sum for any finite source prefix. The formulas are exact even when calculating the integer or traversing the corresponding run is impractical. Simulation-prefix induction, rather than empirical extrapolation, establishes the clock equality.

37.3 Predicted clocks and what was actually executed

For the empty finite tape and standard cofactor, the five-counter prologue from fresh start to the first A0 TM cut takes 142 steps and reaches

$$(L, R, T, H, W) = (0, 0, 0, 15, 0).$$

The exact theorem-derived two-counter cost is

$$79\,936\,151\,060\,302. \tag{24}$$

The saved clock calculation executes the short five-counter trace and sums the prime macro formulas. It does *not* traverse these approximately $8 \cdot 10^{13}$ two-counter steps. A literal trial was stopped at a 5,000,000-step budget; neither a complete two-counter prologue nor a universal computation is reported as executed.

(L, R, T, C)	Predicted two-counter prologue clock
$(0, 0, 0, 1)$	79 936 151 060 302
$(1, 0, 0, 1)$	159 872 302 120 216
$(0, 1, 0, 1)$	239 808 453 180 106
$(0, 0, 0, 13)$	1 039 169 963 779 214

These four cases each have the same 142-step five-counter prologue. Initial scratch $T = 1$, even with $L = R = 0, C = 1$, produces a 459-step five-counter prologue with $H = 47$, and a predicted physical cost exceeding $8 \cdot 10^{40}$. The exact full integer is recorded in `prologue-predicted-clocks.json`. This growth is a concrete reason to separate proof-level simulation from feasible end-to-end testing.

Small macros, history examples and five fresh-entry paths are literally executed by supplementary tests. The independent macro audit checks 2,330 history traces and 27,743 prime traces. These tests help catch implementation mistakes; neither their finite ranges nor their runtime budgets establish universality. The all-input proof comes from complete graph reconstruction, affine identities, invariants and decreasing ranks.

38 The fixed binary five particle instance

Report 15 supplies the independently audited compiler theorem and its explicit finite indexed gate construction [11]. Its source hypothesis is a finite two-natural-counter partial injection with terminal halt and guards constant on finitely many natural-number classes. The present literal

source satisfies those hypotheses on *all* natural IDs by Section 36. The source’s clean-input restriction concerns simulation correctness, and does not weaken the compiler’s global reversibility hypothesis.

The exact byte-preserved compiler proof and authored Report 15 are bundled. Their assertions that their own universal source was not materialized remain correct descriptions of that earlier release. This report provides the new source input separately, without silently editing their history or replacing their finite nonuniversal examples.

38.1 Indexed construction and global reversibility

For context, the compiler assigns an unsigned home mode H_q to every source control and two traveling modes to every moving branch. Every mode has plus and minus versions with distinct gaps in $\{1, \dots, D\}$. A head is a close pair of ones whose gap stores one signed mode. Three isolated ones encode left counter endpoint, origin and right counter endpoint. The initial five-particle support for source ID (q, c_0, c_1) is

$$\text{enc}(q, c_0, c_1)^+ = \{-(Z + c_0), 0, Z + c_1, S, S + d(H_q^+)\}. \quad (25)$$

There are no other nonzero cells, no extra tracks and no varying background.

Each compiler factor is an isolated equal-mass pattern swap. At an eligible key it interchanges two finite supports of the same cardinality that are not translates. Eligibility checks a bounded empty halo, separation from other raw keys, and, when needed, a context predicate identical at the paired endpoints. The isolation lemma proves that the *entire raw-key set* and eligibility are unchanged by the simultaneous rewrite on arbitrary dense or infinite binary inputs. Applying the same factor again restores the input. Each factor is therefore a translation-equivariant local involution and preserves finite numerical mass. A finite composition is globally bijective on the full binary shift; this conclusion does not rely on admissible-input simulation or finite-support injectivity.

On encoded inputs the compiler builds a partial-injective micrograph by subdividing every source move into dispatch, outbound travel, endpoint update, return travel and commit. A zero-update row is one guarded home edge. Intermediate modes carry the original branch identity; domains and exact image predicates make the relevant home joins disjoint. The admissible intermediate states are precisely the subdivisions of *enabled* source edges, not every geometrically plausible traveling packet.

A block E swaps the plus end of each microedge with its minus destination. A block P flips the signed mode. On the admissible doubled graph their factor products realize those matchings. One CA step is $F = P \circ E$. Thus a plus state with successor advances one microedge; a minus state with predecessor moves backward one edge. At a missing edge the sign flips in one step. The global proof is factor-by-factor even on malformed inputs; the matching proof is a separate argument only on the precisely defined admissible set. Report 15 supplies all patterns, contexts and boundary cases.

38.2 All compiler parameters for this source

The defining formulas and their exact values are

Parameter	Definition	Value
m	source controls	122 622
p	moving rows	66 066
a	zero-update rows	75 495
J	natural class cutoff	0
D	$2m + 4p$	509 508
S	$2D + 2$	1 019 018
B_2	$D + 1$	509 509
L_{gate}	$3D + 4$	1 528 528
B_3	$4D + 5$	2 038 037
Z	$10B_3 + 10 + 2J$	20 380 380

The notation L_{gate} distinguishes the compiler isolation constant from the left half-tape integer L . It is called simply L in the machine-readable ledger. The factor families are

$$\begin{aligned}
n_2 &= 4p = 264\,264, \\
n_3 &= 8pD + 23p + m = 269\,290\,886\,364, \\
n_c &= 2p + a = 207\,627.
\end{aligned}$$

Therefore the exact indexed factor count and the unoptimized compositional radius upper bound are

$$N_{\text{factors}} = 8pD + 29p + m + a = 269\,291\,358\,255, \quad (26)$$

$$\begin{aligned}
R_{\text{CA}} &= n_2(6D + 8) + n_3(24D + 32) + n_c(Z + J + 12D + 16) \\
&= 3\,292\,955\,588\,459\,274\,804.
\end{aligned} \quad (27)$$

This is a radius bound, not a claim of least possible radius. The finite local binary truth table could in principle be generated by this indexed composition and would have $2^{2R_{\text{CA}}+1}$ slots at the stated radius. It was not generated. Neither was the list of 269 billion factor objects.

`target-ledger.json` and the loader derive these numbers directly from the pinned source. The original eager `compile_source` was not invoked on this universal table: its quadratic validation and huge object allocation are impractical. A separate source-indexed schema/image validator establishes compatibility without allocating gates. Building an operational lazy compiler would be additional implementation and audit work, not a feature claimed here.

38.3 Loader and exact CA observation clock

The serialized control list puts start first, so $d(H_{\text{START}}^+) = 1$. The standard empty-tape source input is $(\text{START}, 1, 0)$; its CA support is

$$\{-20\,380\,381, 0, 1\,019\,018, 1\,019\,019, 20\,380\,380\}. \quad (28)$$

`five_particle_input` validates the clean source input and the source hash before producing these coordinates. It does not accept a caller's unverified ledger override. Its clean condition $A > 0, B = 0$, no factors 7 or 11, is equivalent to (1), since the 2, 3 and 5 valuations can be absorbed into L, R, T and the remaining factor is coprime to 2310.

For a literal moving branch with update $\delta \in \{-1, 1\}$ and selected pre-counter c , the exact number of CA microedges is

$$\tau_e(c) = 3 + 2(Z + c) + \delta - 4S = 36\,684\,691 + \delta + 2c. \quad (29)$$

A zero-update branch costs one. If the literal source first halts after K primitive rows, let

$$\Theta = \sum_{t=0}^{K-1} \tau_{et}(c_t). \quad (30)$$

The CA first reaches its plus-halt home at time Θ . Reflection *after* that observation is not charged. Executing one local factor is not one CA time step.

[write] Equation (29) is Report 15's clock (5) in Section 8 with this source's constants; Report 18 rewrites it as $\lambda + |(c + \delta)^2 - c^2|$ (equation (54.1)).

The anchored halt word occupies $[0, S + D]$ and has ones exactly at $0, S, S + d(H_{\text{HALT}}^+)$. Its length is $3D + 3 = 1\,528\,527$. The endpoint particles lie outside that interval and distinct signed gaps separate the halt mode from all other modes. This fixed word appears along the valid orbit exactly when the forward source run reaches halt, even though halt is not absorbing.

38.4 Positive return and periodicity now have an unconditional clean family theorem

Corollary 38.1 (Return and periodicity on the explicit family). *For the one fixed binary CA specified above and each clean encoded input x from (1), the following are equivalent:*

1. U halts on the corresponding finite tape;
2. $F^t(x) = x$ for some $t > 0$;
3. x is periodic, or equivalently eventually periodic, under F .

When the input halts with CA clock Θ , the least positive return time and least period are exactly $2\Theta + 2$. Positive return and periodicity are recursively enumerable complete on this computable five-particle input family.

Proof. The source start has no incoming row on any counter values, so its plus home has no predecessor in the unsigned micrograph. A halting run of length Θ microedges follows its distinct plus vertices, flips once at its terminal endpoint, retraces the minus copies, then flips once at the initial endpoint. Partial injectivity and predecessor freedom prohibit repeated unsigned vertices on that finite path: a repetition would propagate backwards to supply an incoming edge at the initial vertex. Plus and minus modes are different, so all $2\Theta + 2$ states before return are distinct. This proves the least-period formula, not just an upper bound.

If U does not halt, Proposition 37.1 supplies an infinite source path. Every enabled subdivision is finite, so the unsigned micrograph path is infinite. It cannot repeat a vertex by the same partial-injectivity and predecessor argument. The CA stays on its distinct plus states forever, never reflecting and never returning. This uses the clean-family no-blocking theorem; a malformed nonfinal stuck source input can have a finite reflected cycle and is outside the claim.

For any bijection, eventual periodicity implies periodicity: from $F^{s+t}x = F^s x$ with $t > 0$, apply F^{-s} to obtain $F^t x = x$. The fixed finite rule can be simulated effectively, so positive return is semidecidable by enumeration of positive times. The universal finite-tape halting reduction provides many-one hardness. This proves the completeness claim. $\square$

Report 15 left this periodicity conclusion conditional on absence of nonfinal blocking. The present source and its clean loader establish that missing premise. The Θ in this corollary uses the *base-source* constants above, rather than raw two-counter step count or a clean-wrapper clock.

38.5 Relationship to the four mass lower bound and exact targets

The companion Report 12 proves exact finite-configuration reachability for fixed one-dimensional conservative CAs with a unique vacuum, positive integer nonvacuum weights and at most one weight-one symbol at mass at most four [10]. A binary number-conserving CA satisfies those model hypotheses. Positive return of x is equivalent to ordinary nonnegative-time reachability of x from $F(x)$; mass conservation keeps the latter input at the same mass. The companion theorem therefore decides positive return at mass at most four. In a reversible CA it also decides eventual periodicity. Together with Corollary 38.1, this gives the sharp threshold five for these questions in the stated binary reversible setting. The lower bound is imported from that byte-preserved companion, not presented as a new independent classification here.

[write] Report 12 is printed as source 17 (Part VI) of `signal-machine-collision-certificates`; its delivered companion copy is not shipped here.

The clean exact-target wrapper of Report 15 also applies because no row enters this source's initial control. Its forward and backward source copies and two bridges have $m' = 245\,245$, $p' = 132\,132$, $a' = 150\,992$, with $J = 0$ after the simple reverse-guard serialization. They return the original input counters to a new terminal control. The resulting exact target depends computably on the input. All compiler constants must be recomputed for this larger wrapper, and its first target clock is $2\Theta' + 2$, its reflected period $4\Theta' + 6$, with Θ' measured using those enlarged constants. That wrapper and its CA remain effective theorem-level constructions here; the present release materializes the base source, not a second giant table.

39 Exact finite horizon quartic certificate counts

Report 15's natural sum-of-squares construction applies to the strict source now actually provided. We state its complete finite formula and specialize its counts, without claiming that the universal-instance polynomial was emitted. Fix a natural input $(q_0, c_{0,0}, c_{0,1})$ and a literal source horizon K . Here K counts *two-counter primitive rows*; it is not a TM-step count, history-macro count or CA time. A new finite polynomial is specified for each fixed K . Varying K outside the polynomial does not create a single unbounded-horizon Diophantine representation.

39.1 Disjoint class normalization and counted witnesses

For $J = 0$, each counter class is exact zero or the tail $\star = \{1, 2, \dots\}$. For every original branch and all four class pairs, retain a normalized branch β exactly when its guard is true on that cell. This partitions the original branch's domain; it does not add instructions to its execution. Let B_{cls} count retained cells, r_{tail} count tail-coordinate occurrences among them, and P_{cls} count retained moving cells. These names avoid confusion with the physical auxiliary counter B and compiler moving-row count p .

The exact enumeration gives

$$B_{\text{cls}} = 350\,054, \quad r_{\text{tail}} = 411\,291, \quad P_{\text{cls}} = 199\,004. \quad (31)$$

They can also be checked by the following class contributions:

Row family	Rows	Cells per row	Tails per row	Moving cells per row
increment	33,436	4	4	4
decrement	32,630	2	3	2
zero test	23,429	2	1	0
positive test	52,036	2	3	0
identity	30	4	4	0

For example a positive guard leaves two cells, with one selected tail in each and one extra tail for the other coordinate in one cell, hence three tail occurrences. Complete cells prevent duplicate witnesses from redundant Boolean formula descriptions.

Give the controls distinct codes $0, \dots, m-1$. Write s_β, t_β for normalized source/target codes, $\delta_{\beta,j}$ for its counter updates, and $\alpha_{\beta,j} \in \{0, \star\}$ for its class. For $K \geq 1$ introduce natural selectors $e_{t,\beta}$ for $0 \leq t < K$, two postcounters at each layer, and one natural slack $z_{t,\beta,j}$ precisely for each tail coordinate. The initial control and counters are external parameters or specialized constants; they are not counted as witness variables.

39.2 Every residual and the full natural zero fiber

Use these integer polynomial residuals, for all indicated indices:

$$A_t = \sum_{\beta} e_{t,\beta} - 1, \quad (32)$$

$$Q_0 = \sum_{\beta} s_{\beta} e_{0,\beta} - q_0, \quad (33)$$

$$Q_t = \sum_{\beta} s_{\beta} e_{t,\beta} - \sum_{\beta} t_{\beta} e_{t-1,\beta}, \quad t > 0, \quad (34)$$

$$U_{t,j} = c_{t+1,j} - c_{t,j} - \sum_{\beta} \delta_{\beta,j} e_{t,\beta}, \quad (35)$$

$$E_{t,\beta,j} = e_{t,\beta} c_{t,j}, \quad \alpha_{\beta,j} = 0, \quad (36)$$

$$T_{t,\beta,j} = e_{t,\beta} (c_{t,j} - 1) - z_{t,\beta,j}, \quad \alpha_{\beta,j} = \star, \quad (37)$$

$$F_K = \sum_{\beta} t_{\beta} e_{K-1,\beta} - h. \quad (38)$$

There is one class residual for every (t, β, j) , either (36) or (37). Let $\mathcal{P}_K$ be the sum of their squares, including all the one-hot, control, update and final residuals. Each residual has degree at most two, so $\deg \mathcal{P}_K \leq 4$. This is a finite ordinary integer polynomial representation, with no floor, exponentiation or unbounded predicate atoms.

Theorem 39.1 (Complete natural witness fiber). *For the fixed natural initial input, the full natural zero fiber of $\mathcal{P}_K$ is empty or a singleton. It is a singleton exactly when the literal source first reaches its designated halt after exactly K rows. The uniqueness includes every selector, postcounter and slack.*

Proof. A sum of real squares vanishes exactly when every residual vanishes. Over $\mathbb{N}$, $A_t = 0$ forces one selector to be one and the others zero. The control residuals make that branch start at the actual current control, and the update residuals impose its exact natural postcounters. A selected exact-zero class forces its old counter to zero by (36). A selected tail forces $z = c_{t,j} - 1 \geq 0$ by (37), exactly the positive class. An inactive tail has $z = 0$, so no unconstrained slack remains. Thus each selected original row is enabled. The terminal residual enforces halt, and no earlier layer can have halted because halt has no outgoing row.

Conversely a legal first-halt trace determines its one-hot cell selectors, postcounters, selected tail slacks, and zero inactive slacks. These satisfy every residual. Deterministic source semantics and the disjoint cell partition make that entire tuple unique by induction on layers. $\square$

At $K = 0$, use $(q_0 - h)^2$ with no witnesses. Its empty tuple is the unique zero exactly when $q_0 = h$. This special case is not obtained by interpreting a negative last-layer index in (38).

39.3 Exact cost and optional clocks or real domain

The core counts for $K \geq 1$ are

$$V_K = K(B_{\text{cls}} + 2 + r_{\text{tail}}) = 761\,347K, \quad (39)$$

$$S_K = K(4 + 2B_{\text{cls}}) + 1 = 700\,112K + 1. \quad (40)$$

Residual slots are counted even if specializing an input makes a row identically zero. Before collecting or substituting large initial constants, the raw residual monomial-slot bound is

$$K(7B_{\text{cls}} + P_{\text{cls}} + r_{\text{tail}} + 5) + 2 = 3\,060\,678K + 2. \quad (41)$$

If residual i has L_i collected terms, expanding the sum of its squares has at most $\sum_i L_i^2$ ordered occurrences before final collection. Thus squaring is not cost-free; these bounds do not claim a small expanded sparse polynomial.

The exact CA clock can be included with one natural variable Θ and one additional residual

$$\Theta - \sum_{t,\beta} e_{t,\beta} \tau_{\text{original}(\beta)}(c_{t,j_\beta}). \quad (42)$$

It is quadratic and uses the unchanged base compiler constants in (29). This adds one variable and one square, and at most $1 + 549\,058K$ raw monomial slots. Its value is uniquely forced. Alternatively Θ can be an external parameter, adding only the square. A deadline N can add one natural slack d and square $(N - \Theta - d)^2$, uniquely fixing the slack when the deadline is met. None of these devices combines all horizons into one uncharged polynomial.

To obtain the same full zero fiber over the nonnegative real orthant, add for every layer

$$\sum_{\beta} e_{t,\beta}^2 - 1. \quad (43)$$

Together with nonnegativity and (32), this forces one-hot selectors: expanding the square of their sum gives a sum of nonnegative pair products equal to zero. The update recurrence from natural initial counters then forces natural postcounters and slacks. This costs K more squares and $K(B_{\text{cls}} + 1)$ raw slots, with no extra variables and degree still at most four. The unlocked orthant count is $700\,113K + 1$ squares. This is not an unrestricted-real theorem; dropping the orthant requirement breaks the one-hot argument.

For input height $M_0 = \max(c_{0,0}, c_{0,1})$, every valid postcounter and slack is at most $M_0 + K$. The largest moving clock constant is $\kappa_{\text{max}} = 36\,684\,692$, so the accepting clock obeys

$$0 \leq \Theta \leq K(36\,684\,692 + 2M_0) + K(K - 1). \quad (44)$$

These are bounds on values at the unique witness, not the complexity of finding it. The input encoder can already make M_0 enormous.

No polynomial with hundreds of thousands of variables per source step is emitted in this release. Report 15 includes fully emitted small example certificates; its generic construction and the present exact source-class ledger justify (39)–(44). The new result here is the universal source instantiation and independently checked counts, not a materialized giant coefficient array or an unbounded single-fold Diophantine claim.

39.4 A smaller certificate using shared primitive test offsets

The full-cell construction is general, but the present source uses only five primitive forms. A specialized construction can avoid its cell copies and per-cell slacks. Let n denote the original branch count, n_Z, n_P the numbers of zero and positive identity tests, and n_M the moving-row count. A decrement is enabled exactly when its pre-step counter is positive; increments and identities are unguarded. This specialization does not accept arbitrary Boolean guards, nonzero thresholds or guarded increments. Its certificate theorem needs determinism and terminal halt, but not source reversibility.

For every step $0 \leq t < K$, use selectors $e_{t,e}$ for the n original rows, shifted postcounter coordinates $z_{t+1,0}, z_{t+1,1}$, and two shared natural offsets $b_{t,0}, b_{t,1}$. Write $\mathcal{P}_j$ for the positive-test rows testing counter j , and define actual counter expressions by

$$C_{0,j} = c_{0,j}, \quad C_{t,j} = z_{t,j} + b_{t-1,j} \quad (t \geq 1). \quad (45)$$

Keep the one-hot, state and final rows (32)–(34) and (38), replacing the cell index β by the original row e . Replace all update/class rows by

$$U_{t,j} = z_{t+1,j} + b_{t,j} - C_{t,j} - \sum_e \delta_{e,j} e_{t,e}, \quad (46)$$

$$B_{t,j} = b_{t,j} - \sum_{e \in \mathcal{P}_j} e_{t,e}, \quad (47)$$

$$G_{t,e} = e_{t,e} z_{t+1,i_e} \quad (e \text{ a zero test on } i_e). \quad (48)$$

The zero guard in (48) deliberately uses a *shifted postcounter*. There are no positive-test slack variables. All residuals have degree at most two; their sum of squares has degree at most four. The offsets are shared by all positive tests on a counter at that step, rather than duplicated per row.

Theorem 39.2 (Primitive shared offset fiber). *For fixed natural initial input and $K \geq 1$, the full natural zero fiber of the shared-offset polynomial is empty or a singleton, and is a singleton exactly for first halt after K original literal rows. The same conclusion holds over the nonnegative real orthant after adding the K selector-norm rows (43).*

Proof. The one-hot rows force one original selector to be one on the natural domain. Equation (47) then fixes each offset to one exactly when the selected row is a positive test on that counter, and zero otherwise. Thus all active and inactive offsets are uniquely determined. Equation (46) imposes the exact update on the actual counters C , which are natural because they are sums of natural coordinates. Check the enabled domain by primitive type:

- An increment or unguarded identity has no guard to prove
- For a selected decrement, both current offsets vanish; its natural shifted postcounter is the old selected counter minus one, forcing that old counter to be positive
- For a selected positive test on j , the update is zero and $b_{t,j} = 1$, so its old and new actual counter is $z_{t+1,j} + 1 \geq 1$
- For a selected zero test on j , the update and both offsets are zero; (48) gives $z_{t+1,j} = 0$, hence its old actual counter is zero

The state and final rows therefore give a legal trace ending at halt, with no earlier halt possible because it has no exit. Conversely a legal first-halt trace fixes its selectors, actual postcounters and offsets, then uniquely fixes $z_{t+1,j} = C_{t+1,j} - b_{t,j}$. This is nonnegative for a selected positive

test by its guard, and is just the natural postcounter otherwise. Every residual vanishes and no unused coordinate is free.

For nonnegative real witnesses, the added norm rows and one-hot rows force one-hot selectors by the same pair-product argument used earlier. The updates from natural initial constants and the offset equations then force integral actual counters and integral shifted coordinates. The remaining proof is unchanged. No unrestricted-real claim follows. $\square$

At $K = 0$ retain the separate constant residual $q_0 - h$ and no core witnesses. A clock, when requested, contributes its single variable and residual even in this special case, forcing clock zero. A stuck nonhalt never supplies a witness.

39.5 The improved counts and the price of expansion

There are $n + 4$ variables per layer. Each layer has one one-hot row, one state row, two update rows, two offset rows, and n_Z zero-guard rows. With one final row this gives

$$V_K^{\text{prim}} = K(n + 4), \quad S_K^{\text{prim}} = K(6 + n_Z) + 1. \quad (49)$$

The explicit written residual-term counts, retaining zero coefficients before collection, are:

Row family	Written term slots
one-hot	$K(n + 1)$
state	$n + 1 + (K - 1)2n$
updates together	$n_M + 6 + (K - 1)(n_M + 8)$
offsets	$K(n_P + 2)$
zero guards	Kn_Z
terminal	$n + 1$

Their exact sum is $K(3n + n_M + n_P + n_Z + 11)$. The initial constants count as atomic terms. The first-layer update saving of two terms cancels the state/final boundary contribution of two.

For the literal universal table, $n = 141\,561$, $n_Z = 23\,429$, $n_P = 52\,036$ and $n_M = 66\,066$. Thus the improved core has

$$V_K^{\text{prim}} = 141\,565K, \quad S_K^{\text{prim}} = 23\,435K + 1, \quad (50)$$

$$L_K^{\text{written}} = 566\,225K. \quad (51)$$

The generic counts (39)–(40) remain correct for their broader interface. These are two different constructions; the improved count does not silently remove coordinates from the old polynomial.

The clock row is (42) with old counter $C_{t,j}$ from (45), so remains quadratic. It adds one variable and one square, and exactly

$$1 + K(n + 2n_M) - n_M = 273\,693K - 66\,065 \quad (52)$$

written terms for $K \geq 1$. Its first-layer old counter is one constant instead of a sum of two coordinates. The paid real norm adds K squares and $141\,562K$ written terms, but no variables. All witness-height bounds above remain valid; offsets and selectors are at most one, and shifted postcounters are at most the corresponding actual counters.

One can substitute the selector sums from (47) and remove the two offset variables and equations. That reduces the core to $141\,563K$ witnesses and $23\,433K + 1$ squares, but increases its written

terms to $618\,255K - 52\,034$. More importantly, the old counter in each later moving clock then contains a sum of all previous positive-test selectors on that counter. Since

$$\sum_{j=0}^1 n_{M,j} n_{P,j} = 2\,342\,126\,148,$$

the substituted clock costs $2\,342\,333\,775K - 2\,342\,126\,147$ written terms. The two shared offsets preserve a much smaller clock expression. These are exact costs of specified formulas, not optimality lower bounds.

The indexed implementation computes exact collected row lengths L_i and $\sum_i L_i^2$ without enumerating the horizon's residual array. This latter quantity charges all ordered products in expanding the squares; it is not the final number of distinct collected monomials. Source snapshots and adjacency maps are stored, while variable and residual indices are calculated from the fixed horizon. Explicit row, full polynomial and witness export have finite size guards. The universal instance is a finite indexed family, not a fully emitted giant polynomial.

The small five-row example in `primitive-certificates/` does provide a complete emitted certificate: 46 variables and 42 squares for source horizon five, exact clock and paid real norms. Its expanded quartic has 1,100 collected monomials, and its entire witness evaluates to zero with clock 3,115. The same five-row source was also compiled into its actual small binary CA: 387 factors, radius bound 199,300 and observer length 63. All 6,232 full CA steps through its first return were executed, with mass five throughout, first plus-halt at 3,115 and no earlier return. Its complete raw orbit and checked receipt are included; this is a finite test machine, not the universal run. The universal horizon-one ledger is separate and uses the clean empty-tape source input (START, 1, 0); it is not an accepting witness or a full universal polynomial.

39.6 Certificates for least period and return multiples

For a clean halting input, Corollary 38.1 gives least period $\Pi = 2\Theta + 2$. With the exact clock retained, a specified natural time T is its least period exactly when one more residual

$$T - 2\Theta - 2 \tag{53}$$

vanishes. This adds one square and no witness, and preserves both the natural and paid real-orthant complete-fiber theorems.

To certify a positive return time instead, append a natural witness u and one residual

$$T - (u + 1)(2\Theta + 2). \tag{54}$$

This costs one variable and one square, stays quartic, and has the unique extension $u = T/\Pi - 1$ exactly when T is a positive multiple of Π ; on such a zero, $0 \leq u < T$. The augmented coordinate-height bound is the maximum of the core-with-clock bound and T . Its fully written expansion has at most five monomial slots. The multiplier must remain *natural*, or be explicitly a nonnegative integer coordinate in a mixed-domain statement. Allowing all witnesses to be nonnegative real would incorrectly accept $T = \Pi + 1$ with $u = 1/\Pi$; the selector-norm proof does not make this newly introduced coordinate integral. These two period rows are theorem-level additions, not advertised CLI options of the bundled module.

There is a standard structural obstruction to replacing that integer multiplier by finitely many purely real polynomial witnesses. Fix one halting input and its horizon, hence a period $\Pi \geq 2$. A projection of a finite real polynomial equation/inequality system is semialgebraic by Tarski–Seidenberg [9, Theorem 2.1]. In one variable its defining polynomial signs stabilize beyond their finitely many roots, so membership on sufficiently large natural T is constant. Positive multiples

of Π and their complement are both unbounded. Therefore no such fixed finite real-witness system represents exactly those natural return times, even without a uniqueness requirement. This does not rule out (53), mixed integer witnesses, or fixed finite bounded-time formulations.

40 Evidence and portable reproduction

40.1 Independent layers of evidence

The evidence has several distinct scopes. Mathematical induction supplies unbounded quantifiers. Whole-graph reconstruction connects the proof graphs to the saved rows. Executable affine paths check loop bodies on arbitrary natural parameters. Finite interpreter runs provide additional implementation checks. The claims should not be interchanged.

Check	Exact scope
Primary transcription comparison	All 30 TM cells; the recorded visual review also checks the initial scanned symbol
Virtual affine checker	437 affine paths covering all 528 combined instructions and all 29 defined TM transitions
Reversible affine checker	54,084 affine paths covering every one of the 5,451 five-counter and 141,561 two-counter rows
Independent reconstruction	Complete edge multisets of splitting, normalization, history, prime expansion and signed serialization
Independent macro interpreter	2,330 history runs and 27,743 prime runs, including clocks and disabled singleton behavior
Supplementary concrete interpreter	444 prime examples, 2,330 history examples and five literal fresh-entry paths
Separate release audit	2,349 small TM macros; all-natural symbolic source/image masks; API rejection and isolation; five prologue predictions; regeneration; normal and optimized mutation rejection
Primitive certificate checks	Thirteen exact implementation tests in both modes; independent 53,786-check audit with 136 expected rejections per mode; source ledger and bounded-allocation checks; complete small polynomial and its 6,232-step actual CA orbit

The supplementary concrete counts include 53,990 prime primitive steps, 58,250 history primitive steps and 350 fresh-entry primitive steps. They do not include a traversed full physical prologue. The new primitive-certificate example additionally executes all 6,232 actual CA steps to its first return, with the clock and period given in Section 39.4. The larger independent prime run count is a separate test suite and is not substituted for the affine coverage count.

The independent release-level audit reviewed the all-input composition and periodicity argument, reproduced the class ledger and CA constants, and checked byte identity of the relevant previous front-end and compiler snapshots. It is explicit about inherited dependencies: neither that audit nor the saved transcription comparator re-proves the Neary–Woods universality theorem from scratch, repeats the primary visual inspection, or independently proves Report 12’s four-mass theorem.

40.2 Checker hardening before the final freeze

An audit found that early checker versions used Python assertions for proof obligations. A deliberately corrupted serialization was rejected under normal Python but could produce a passing receipt under `python -O`, which removes assertions. All substantive conditions were replaced by explicit runtime checks and raises. The final eight stages pass in normal and optimized Python, and the independent release audit rejects six controlled corruptions in both modes. These include a serialized row, a prime primitive, the TM table, the virtual program, a prime certificate, and a malformed guard submitted to the class-count helper.

The class helper also pins the exact source and checks its $J = 0$ primitive guard vocabulary before deriving its counts. These fixes concern checker robustness. The mathematical source table and its SHA-256 did not change. The packaged files are the final hardened versions, not the vulnerable intermediate versions. Deep or cyclic arbitrary in-memory guard objects are not part of the finite parsed JSON source interface; no broad hostile-object parser-security theorem is asserted.

40.3 Release inventory and exact manifest

The portable release contains:

- This PDF, its complete LaTeX sources and a rebuild script
- The literal source, all saved intermediate counter tables, the pinned TM/front-end dependency, loader, generator, proof notes, eight substantive checkers and deterministic receipts
- The independent macro and periodicity proofs, plus the finalized release-level audit script, two mode receipts and stabilized hash record
- One unchanged compiler implementation/schema/proof snapshot and the byte-preserved authored Report 15 and Report 12 companion documents with their required LaTeX sections and vector figures
- A standard-library replay driver and an exact per-file size/SHA-256 manifest for this release
- The specialized indexed primitive certificate module, complete proof, normal/optimized test receipts, one fully expanded small example with its entire witness, and that example's complete actual CA orbit

Third-party paper PDFs, screenshots, working logs, cache directories, private scratch replay trees and previous huge release archives are not included. The source's old package-wide manifest is not reused after those exclusions; the present manifest covers exactly the files actually shipped. The audit scripts use package-relative defaults and mandatory bundled dependency pins. Their receipts were regenerated for this portable tree; `PACKAGING-ADAPTATIONS.md` records the exact script and metadata adaptations. The current manifest binds the final document and release bytes.

40.4 One fresh standard library replay

From a fresh extraction, run `python3 replay.py`. It first verifies the release inventory, lengths and SHA-256 hashes. It performs the substantive source checks and regeneration inside a temporary copy, so the shipped files and historical receipts remain unchanged. It invokes the independent release-level script in normal and optimized modes on private copies. It also runs the specialized certificate tests and source audit in both modes, regenerates the small polynomial artifacts, and

replays every step of the small certificate example’s CA orbit. The final result is a concise passing receipt written outside the published tree when requested. No network access or third-party Python package is needed for these checks; Python 3.10 or newer is required. Rebuilding the PDF separately uses `sh build.sh` and a normal TeX Live installation.

The generator must reproduce every byte of seven generated JSON files:

- `primitive3.json`, `normalized3.json`, `reversible5.json`
- `reversible2-primitives.json`, `source.json`
- `certificates.json`, `build-stats.json`

The driver verifies those hashes after regeneration. It never invokes the eager universal CA compiler, allocates a giant factor list, simulates a full physical prologue, or claims to re-run the earlier companions’ entire computational suites. Their byte-preserved proofs and prior test descriptions are included as bounded dependencies.

The command `python3 source/loader.py --left 101 --right 01 --particles` demonstrates the finite loader. Its huge-integer computation can become expensive for long tape halves; the exact API and clean domain are the relevant guarantee, not a uniform operational budget.

40.5 What remains unmaterialized

The source table is a 32,034,272-byte expanded JSON document. The alternative primitive serialization and every intervening source table are also present. This should not be confused with execution scale: even the empty tape’s initialization is immense after prime coding. The following are deliberately outside the delivered materialization claim:

1. A full physical universal run or complete two-counter initialization
2. The CA’s approximately 269-billion-element factor array or its local binary truth table
3. A universal-instance expanded quartic polynomial or full accepting universal witness
4. A proof-assistant formalization, optimized machine/radius, practical universal emulator, or new literature priority result

These boundaries leave a precise positive result: one exact finite literal source with proved all-natural partial injection, a complete clean-input universality reduction, state-dependent exact clocks, and an auditable interface to the already specified finite binary CA.

41 Further questions that follow from this release

The most useful next work is to reduce a specific bottleneck or strengthen a particular interface while preserving the exact proof obligations.

1. **Reduce the initialization cost.** The blank-input history of 15 already creates a large power of 7. Can a different incoming-edge ordering, source normalization, or reversible prologue reduce that history without introducing overlapping ranges? Compare exact prologue clocks, not only control counts, and preserve the no-incoming start normal form. [write] Answered by Part III: Report 17’s four history-tag swaps reduce empty startup from the derived 79,936,151,060,302 source steps to 138 executed ones, keep the no-incoming start, and are optimal within the fixed templates (Theorem 47.1); Report 18 proves the same comparison for the automaton’s clock.

2. **Build a lazy indexed compiler.** Implement random access to the finite gate families and a verified on-demand local evaluator, avoiding the eager object array. Audit boundary cases, actual image guards and factor order on small fully materialized instances before applying its indexing to the universal source. A lazy implementation would improve access to the specification; it would not by itself make the enormous radius efficient. [write] Answered by Part IV: Report 19’s lazy evaluator (Theorem 65.1), with arithmetic factor access and no factor array.
3. **Study further certificate tradeoffs.** The shared-offset construction already improves the general class partition for this primitive source. Eliminating its two wires saves coordinates but makes the clock’s written expansion much larger. Can additional structured sharing reduce the remaining selectors or expanded square cost while preserving degree four, exact clocks and complete-fiber uniqueness? Any comparison must include every wire, equation, coefficient and expansion cost, rather than counting witness coordinates alone. [write] (3 October 2026.) Reports 20 and 28 (Parts V and VI) give further fixed-horizon quartic certificates for the automaton itself, both with fully charged ledgers: Report 20 for the ordered rule under an event budget (Section 84; only a mass-two exporter), Report 28 for the new parallel rule, with exact syntax, coefficient and expansion counts (Section 118). Neither revisits this Report’s shared-offset certificate, so the question stays open as stated.
4. **Mechanize the small trusted core.** Formalize natural-counter primitive semantics, the source/image disjointness criterion, macro loop invariants and the composition lemma. Then connect a verified JSON parser and graph reconstruction to the pinned table. The current checks and independent reviews are not a substitute for that formalization.
5. **Separate source size from physical complexity.** Which exact tradeoffs between controls, history growth, prime ordering, factor count and radius are attainable? A smaller source need not produce a faster first observation, and reducing a summed radius bound need not alter the true local dependency radius by the same amount.
6. **Strengthen execution evidence at an honest scale.** Find and publish short source instances whose complete literal and CA traces fit a reasonable budget, including both halting and indefinitely extendable test families. Accelerated macro traces should carry verified clock summaries and be labeled as accelerated. They must not be described as raw execution of the astronomical primitive trajectory.
7. **Investigate stronger algebraic goals separately.** The exact finite-horizon fibers here do not resolve an unbounded single-fold Diophantine representation problem. Any proposal combining all source horizons must explicitly account for that quantification, polynomial size, auxiliary witnesses and their uniqueness.

The immediate gap from Report 15 is closed: universality no longer rests on an uninstantiated reversible-source existence citation. The finite table and its clean loader are now fixed, auditable data. The remaining large objects have exact finite specifications and counts, with their nonmaterialization and practical limits stated explicitly.

Part III

Startup histories and cellular clocks (Reports 17 and 18)

Startup history in a literal reversible source

Four tag swaps preserve every clean input and give earlier literal arrival at every TM cut

Report 17 Separate construction and reproducible audit

3 October 2026

Abstract of Report 17, as delivered. Four local swaps of incoming history tags remove the empty-startup explosion in the literal reversible universal two-counter source of Report 16. The new fixed source has the same 122,622 controls and 141,561 expanded rows. It preserves global partial injection on all natural configurations and the entire original clean-loader domain, including nonzero initial scratch. For scratch $T = 0$, the exact startup clock is $76A + 4\lfloor A/5 \rfloor + 24\lfloor A/7 \rfloor + 48\lfloor A/11 \rfloor + 62$, where $A = C2^L3^R$. Empty startup was actually traversed in 138 serialized steps, followed by exactly one TM transition in 210 more steps. The old 79,936,151,060,302-step startup is a theorem-derived comparison, not an executed run. An order-preserving embedding of primitive operations proves that runs from the same clean input reach every corresponding normalized boundary no later, and every TM cut strictly earlier, in the new literal machine. This is compatible with an executed 28-to-104-step slowdown from an unrelated equal-history internal state. Within the fixed recorder and prime templates, a first-encounter theorem also proves that the six startup choices are optimal. The report gives the unbounded proof, exact delta, nonzero-scratch formula, trace accounting, and a pinned offline replay package. No raw cellular-automaton speed comparison or universal-program completion is claimed.

Release separation. This is a new machine and a new indexed CA instance. Report 16 and its frozen artifacts are not revised by this report. Equal counts and geometry do not identify the two CA maps.

Evidence levels. The mathematical statements below are ordinary proofs supplemented by independent graph checks and finite executions. They are not proof-assistant formalizations, novelty claims, efficiency claims for universal computation, or machine-size records.

[write] Delivered as `Reversible_Startup_Optimization_Package.zip` (batch 82, archive 16; arrival `db37d18c8`), dated 3 October 2026; files placed by `bca6383e9` with the prefix given in Appendix A. Report 17 starts from Report 16's frozen source (Part II) and answers its first question.

42 The machine and the precise comparison

Let M_{old} be the frozen literal machine of Report 16 and let M_{new} be the present `source.json`. Every branch is one increment, decrement, zero test, positive test, or identity on exactly two natural counters. A successful test counts as a step. No history macro, prime operation, Python callback, or hidden register is available as a single instruction to either literal machine.

The simulated data are three natural counters (L, R, T) . The five-counter history layer adds (H, W) ; the literal two-counter layer encodes a boundary by

$$N = C2^L3^R5^T7^H11^W, \quad B = 0, \quad C \geq 1, \quad \gcd(C, 2310) = 1. \quad (42.1)$$

The *clean loader* means $H = W = 0$ initially. It does not require $T = 0$. A finite binary tape with scanned symbol zero is loaded by taking L and R to be its nearest-head-first half-tape binary integers and using $T = 0$, $C = 1$. Initial data are finite and computable, though possibly enormous.

Theorem 42.1 (Preserved semantics and improved paired clocks). *For every $L, R, T \in \mathbb{N}$ and positive C coprime to 2310, compare the two runs from the same physical input*

$$(\text{START}, C2^L 3^R 5^T, 0).$$

They simulate the same normalized three-counter path and hence the same Neary–Woods universal TM computation. Each enabled simulated edge completes in finite time. Each machine reaches HALT if and only if that source computation halts. Both literal step relations are partial injections on all natural configurations, including unencoded and unreachable ones.

At every corresponding normalized-source boundary, the cumulative literal two-counter clock of M_{new} is no larger. At every TM cut, including the first cut after initialization, it is strictly smaller. These assertions hold for every finite prefix of an infinite run as well as a halting run.

A normalized boundary means a completed history macro (or a completed nonrecording edge), not an arbitrary internal physical control. A TM cut also requires $T = 0$; a pop loop can revisit a cut label with nonzero scratch. The theorem compares source-step clocks. It supplies neither an ordering of all intermediate physical states nor a comparison of raw CA evaluation time.

The new literal source SHA-256 is

fa61d06178d718511232a80a05e02192d940fde3c273f0c202f0623253b2ade3

The frozen comparison source SHA-256 is

38c586706fa1069442d3e5103151adb156448b7d7c783f5c46d9d507b87fe73a

These identify different byte strings. The unchanged lower three-counter inputs, complete reconstructed graphs, and a mandatory pinned comparison baseline make the difference reproducible without consulting another release.

43 The exact four swaps

The normalized three-counter source has indegree at most two. At each of its 233 colliding incoming pairs, reversibilization chooses a bijection to $\{0, 1\}$. The choice of which edge gets zero is not a semantic property of that data operation. The new policy assigns zero to six distinct startup edges. Two already had zero; four pairs are actually permuted.

Destination	Startup edge	Old	New	Companion edge
init_clear_T	entry	0	0	v0000d
n0001_0	v0000z	0	0	v0003
n0001_1	n0001e0	1	0	v0028
n0001_2	n0001e1	1	0	v0066
n0001_3	n0001e2	1	0	v0155
tm_A0_pop0	n0001e3	1	0	v0303z

Table 6: The six explicit policy choices. Each companion has the opposite bit. The last four rows are the complete history-label delta.

No pair is asked to give both edges zero. All other 227 pairs retain their original order; the first two policy pairs also retain theirs. Thus 229 of 233 pair assignments are unchanged. The first three changed companions are counter- L increments from `tm_A1_0_write`, `tm_B1_0_write`, and `tm_E0_0_write`. The fourth is the $T = 0$ test from `tm_I0_0_restore`.

The primitive and normalized three-counter files are byte-identical to the predecessor. Splitting combined instructions, inserting the fresh start, normalizing indegrees, selecting primes, and expanding prime graphs all retain the same rules. The different tag assignment changes the expanded reversible rows and therefore the fixed literal machine, while preserving all numerical dimensions.

Layer	Controls	Rows
Split three counters with fresh start	763	995
Normalized three counters	792	1,024
Reversible five counters	4,520	5,451
Literal reversible two counters	122,622	141,561

Table 7: All four counts agree exactly with the predecessor.

At the literal layer the census remains 33,436 increments, 32,630 decrements, 23,429 zero tests, 52,036 positive tests, and 30 identities. The five-counter control and row-name sets are unchanged, with exactly eight named rows changed. At the expanded two-counter layer, enumeration propagates the swaps into private-control names; equality of the old and new control-name sets is not claimed. The label change is not a row-deletion argument, a scratch restriction, or a reduced-state construction.

The exact expanded census replaces 51 control names by 51 new names and 63 row names by 63 new names. Among rows sharing a name, eight differ; 622 row positions differ when the serialized lists are compared in order. These are representation-level delta measures, not 622 new history choices. The full old/new comparison is in `exact-delta-receipt.json`.

44 Why the whole original input theorem survives

44.1 The history invariant is independent of the bit assignment

A recorded edge first performs its original data operation. With boundary history $H = h$ and work $W = 0$, the transfer loop then has invariant

$$(H, W) = (h - j, j), \quad 0 \leq j \leq h.$$

One transfer iteration is the four-operation word

$$P_H, -H, +W, P_W.$$

Its decreasing rank is H . The exit is Z_H . Bit one then inserts $+H, P_H$; bit zero inserts nothing. The doubling loop has invariant

$$(H, W) = (b + 2j, h - j), \quad 0 \leq j \leq h,$$

and six-operation word

$$P_W, -W, +H, P_H, +H, P_H.$$

Its rank is W , and its exit is Z_W . Every decrement and following positive test is enabled by the invariant. The original data operation's domain is preserved because it is performed before this history work.

Lemma 44.1 (History recording). *For either $b \in \{0, 1\}$, every enabled recorded edge terminates from $(data, h, 0)$ at its original destination with the exact post-operation data and $(H, W) = (2h + b, 0)$. Its five-counter duration is*

$$1 + 1 + 4h + 1 + 2b + 6h + 1 = 10h + 4 + 2b. \quad (44.1)$$

A nonrecording edge costs one five-counter step and leaves H, W unchanged.

Proof. The two displayed invariants and their strictly decreasing ranks prove termination and the output. The seven summands count the source operation, entry Z_W , transfer body, transfer exit, preparation, doubling body, and final exit. No part assumes a prescribed ordering of a collision pair. $\square$

Zero history can remain zero after appending a zero bit. This does not make the machine irreversible: the companion input reaches the same destination with odd history, and internal joins retain disjoint tests. The complete configuration, including finite control, determines the predecessor. The history integer alone is not asserted to encode every arbitrary bit string with its length.

44.2 Prime graphs and exact literal costs

The unchanged prime encoding (42.1) represents the five counters using 2, 3, 5, 7, 11. For an enabled five-counter primitive at positive encoded integer N , acting on prime p , its expanded two-counter cost is

primitive	literal two-counter steps	
0	1	
+	$(p + 7)N + 3$	(44.2)
−	$4N + (p + 3)(N/p) + 3$	
Z or P	$4N + 4 \lfloor N/p \rfloor + 3.$	

These are costs of explicitly saved paths, not accelerated instructions. Unique factorization identifies a zero exponent with nondivisibility and a positive exponent with divisibility. A test records a bounded remainder and restores the original N using a destination-shared restorer. Multiplication and division consume a finite work counter; division is enabled only on a whole number of prime-sized groups. The complete invariants and decreasing ranks are bundled in `independent-macro-audit.md` and `PROOF.md`. Private controls are separated from boundaries and from each other, except for explicitly shared restorers. Per-edge independent restorers would not provide the required disjoint incoming images.

All costs in (44.2) are finite and positive, and nondecreasing in N on their enabled domains. This monotonicity will be used only after an operation-by-operation matching has been constructed.

44.3 Global injection and nonblocking simulation are separate

Each literal primitive is individually injective on its domain. In both reversible graphs, outgoing multiple edges have complementary zero/positive domains on the same counter; incoming multiple edges have complementary zero/positive images on the same counter. Complete graph reconstruction and all such joins are checked again after relabeling. Hence every natural configuration has at most one successor and at most one predecessor.

The serialized-source checker independently derives image guards. A decrement's image is all naturals; an increment's image is positive; a test's image is its guard. Since the class cut is $J = 0$, every predicate is constant on each of the four product classes obtained from zero/positive in

the two counters. Exhaustive checks of those classes establish the all-natural property, not a finite-grid extrapolation. The start has no incoming row and the halt has no outgoing row.

For a promised clean input, the normalized data source always has an enabled nonhalt continuation. Lemma 44.1 and the prime ranks expand it in finite positive time without a premature halt. Induction on source edges yields the same finite data prefixes and halting equivalence. An infinite source run produces infinitely many completed finite macros, so cannot become stuck or spend infinitely long inside one macro. This proves the semantic assertions of Theorem 42.1. It does not promise successful simulation from arbitrary dirty-work inputs, even though global partial injection holds there too.

45 Exact initialization for all natural scratch values

45.1 Zero scratch keeps the entire history at zero

For $T = 0$, the six startup bits are now 0,0,0,0,0,0. Every history macro therefore has $h = 0, b = 0$, costs four five-counter steps, and preserves $H = W = 0$. The entire startup consists of exactly five identities, one Z_T , six Z_H , and twelve Z_W . Neither L nor R changes.

Set $A = C2^L3^R$. At every prime boundary of this prologue the encoded integer is A and the auxiliary counter is zero. Define $F_p(A) = 4A + 4\lfloor A/p \rfloor + 3$. Summing the 24 actual five-counter operations gives the following exact result.

Proposition 45.1 (Zero-scratch startup clock). *For all $L, R \in \mathbb{N}$ and admissible C , the startup from (START, $A, 0$) to the first tm_AO_pop0 boundary has duration*

$$\begin{aligned} P_0(A) &= 5 + F_5(A) + 6F_7(A) + 12F_{11}(A) \\ &= 76A + 4\lfloor A/5 \rfloor + 24\lfloor A/7 \rfloor + 48\lfloor A/11 \rfloor + 62. \end{aligned} \quad (45.1)$$

It ends with five-counter data $(L, R, 0, 0, 0)$ and physical pair $(A, 0)$.




Figure 5: History at the six clean $T = 0$ startup collisions, starting at history zero. Edge labels are appended bits; each update is $H \mapsto 2H + b$. The numerical physical encoding amplifies the final old history to a factor 7^{15} . The drawing is a boundary diagram, not a physical trajectory.

For $A = 1$, (45.1) gives 138. The old boundary history was 15, old five-counter startup took 142 steps, and its exact predicted literal clock was 79,936,151,060,302. That old full physical run was not traversed. Both the new formula and the old comparison are derived from finite macro words and exact prime costs; only the new small literal startup is also backed by direct full execution.

45.2 No nonzero-scratch input is removed

Let $T \geq 0$ and $h_T = 2^T - 1$. The initial entry records zero. Each of the T successful scratch decrements records one, preceded by its nonrecording positive test. After i clears the history is $2^i - 1$. The zero exit and four merge edges then append five zeros. Therefore

$$H_{\text{new}} = 32(2^T - 1), \quad H_{\text{old}} = 32(2^T - 1) + 15, \quad W = 0. \quad (45.2)$$

The data at the end are $(L, R, 0)$ in both machines.

The exact new five-counter clock is

$$\begin{aligned} P_5(T) &= 4 + \sum_{i=0}^{T-1} (10(2^i - 1) + 7) + \sum_{j=0}^4 (10 \cdot 2^j h_T + 4) \\ &= 320 \cdot 2^T - 296 - 3T. \end{aligned} \tag{45.3}$$

For $T = 0$, the empty first sum gives 24. The old five-counter clock is exactly $P_5(T) + 118$ for every natural T . One way to see the constant 118 is to compare the five final macros: their old added-history increments are 0, 1, 3, 7 before the last four relevant costs, contributing $10(0 + 1 + 3 + 7)$, plus four two-step preparations, for $110 + 8$.

At $T = 1, L = R = 0, C = 1$, the new five-counter startup takes 341 steps and ends with $H = 32$. Its composed literal clock is

$$18,595,602,803,420,561,611,627,873,065.$$

This is a theorem-predicted value, not a traversed two-counter run. Formula (45.1) must not be applied to this input. Nonzero scratch is fully supported semantically, but its recorded clearing history remains in an exponent and can be prohibitive.

46 The order preserving comparison proof

A smaller final history by itself does not prove a smaller runtime: the intermediate work counter also enters the prime encoding. The needed comparison matches actual primitive operations in chronological order.

Lemma 46.1 (Recorded macro domination). *Fix equal data and cofactor and an enabled normalized recording edge. Let its old history be $h + d$ and new history h , with $h \geq 0, d \geq 1$, and let the old and new tags be any $b_o, b_n \in \{0, 1\}$. Its full old literal two-counter duration is strictly larger than its new duration.*

Proof. Match the common data operation, entry Z_W , and first h transfer iterations, preserving order. Each matched pair uses the same symbol and counter, and the old encoded integer is 7^d times the new one. Then allow the old run to perform its d unmatched transfer iterations. Match the two Z_H exits; the ratio is now 11^d .

Three bit cases are immediate. If both tags are one, match the two preparation operations at ratio 11^d . If both are zero, there are no preparations. If $b_o = 1, b_n = 0$, leave the old preparation unmatched. In each of these cases, let the old run complete d surplus doubling iterations. The two work counters now both equal h , and the old history exceeds the new by $2d + b_o - b_n \geq 1$. Match the remaining h doubling iterations and final Z_W ; every encoding ratio is $7^{2d+b_o-b_n} \geq 7$.

The case $b_o = 0, b_n = 1$ needs a different alignment. In the old run's first doubling iteration, leave its $P_W, -W$ unmatched. Match the new preparation $+H, P_H$ to the old iteration's *first* $+H, P_H$ pair. Just before these matches the old work exceeds the new by $d - 1$, so the encoding ratio is $11^{d-1} \geq 1$. Finish that old iteration (its second $+H, P_H$ pair), then complete the next $d - 1$ old doubling iterations. Old and new work now both equal h , while the old history is $2d$ and new history is 1. Match the new h doubling iterations and exit to the old remainder at ratio $7^{2d-1} \geq 7$.

In every case, this is an order-preserving injection from new operations into old operations. No old operation is used twice. Every matched pair has the same prime and primitive, with old

encoded integer at least new. The exact costs (44.2) are nondecreasing on those domains. There are unmatched old positive-cost operations (in particular the extra transfer iterations). Summing proves strict inequality. $\square$

b_o	b_n	Matching before the common doubling suffix	Final ratio
0	0	Skip d old doubling iterations	7^{2d}
1	1	Match preparation, then skip d old iterations	7^{2d}
1	0	Skip old preparation and d old iterations	7^{2d+1}
0	1	Match new preparation inside the first old iteration; finish d old iterations in total	7^{2d-1}

Table 8: The four bit cases in the chronological matching. In the last row, preparation is matched at ratio 11^{d-1} before the common suffix.

Lemma 46.2 (The first strict gain). *At equal incoming history, equal tags give equal costs. At equal history with $b_o = 1, b_n = 0$, the old recorded macro is strictly longer.*

Proof. Match the data operation, entry, transfer, and transfer exit. Skip the old two preparation operations. The matched doubling suffix has the same work but old history larger by one, so its encodings are larger by a factor of seven. Monotonicity plus the skipped positive-cost operations proves the strict claim, also when the doubling suffix is empty. $\square$

Proof of the clock assertions in Theorem 42.1. Both runs follow the same normalized data path. Throughout the common scratch-clearing portion their histories and tags agree. The first changed startup tag is old one/new zero, so Lemma 46.2 yields the first strict gain. Subsequent startup changes also have this favorable sign. At the first TM cut, (45.2) gives a history difference $d = 15$ for every original loader input.

Afterward, a nonrecording edge leaves d unchanged. On a recording edge,

$$d' = 2d + b_o - b_n \geq 2d - 1 \geq d \geq 15. \quad (46.1)$$

Thus the positive gap persists, whatever the later bit choices. Each recorded macro is strictly improved by Lemma 46.1. For a nonrecording edge, old and new use the same operation on equal data, and old encoding is 7^d times new. Prime-cost monotonicity gives a nonlarger new duration (an identity can tie). Induction on normalized edges proves all cumulative inequalities. The strict gain already obtained in startup persists to every TM cut. $\square$

For an arithmetic cross-check, after the data operation put $K = C2^L3^R5^T$. The exact costs of individual loop iterations are

$$\begin{aligned} \text{transfer at } (H, W) = (h - j, j) : & \quad 136K7^{h-j-1}11^j + 12, \\ \text{doubling at } (H, W) = (b + 2j, h - j) : & \quad 474K7^{b+2j}11^{h-j-1} + 18, \\ \text{bit-one preparation after transfer} : & \quad 46K11^h + 6. \end{aligned}$$

Here a loop body is used only within its enabled iteration range. These formulas are obtained by summing four, six, and two corresponding primitive costs. They corroborate the embedding but do not replace its chronological argument. The independent checker also tests 336 actual-row suffix embeddings, covering all four bit pairs and several history gaps and cofactors.

47 The exact optimum within the orientation class

There is a useful stronger conclusion, with a deliberately narrow comparison class. Keep the normalized data program, recorder, prime encoding, and every macro template fixed. An *orientation* chooses one of the two bijections to $\{0, 1\}$ at each incoming recording pair. For this source there are 2^{233} static orientations. This class changes labels only; it does not vary algorithms, registers, primes, normalizations, or universal machines.

Fix a finite enabled normalized trace $e_1, \dots, e_n$, including its data values, and a common initial history $h_0 \geq 0$, $W = 0$, and admissible cofactor. Define a greedy orientation G : when an incoming edge first traverses a previously unseen recording pair, assign that edge zero and its companion one. Choices at pairs never traversed in the trace are arbitrary. Merely starting at a pair's target does not count as encountering the pair.

Theorem 47.1 (Finite-trace first-encounter optimality). *Within the fixed orientation class, G simultaneously minimizes the history and the cumulative literal two-counter arrival clock at every normalized boundary of the finite trace. If another orientation O differs on any encountered pair, then from its first such encounter onward its history and cumulative clock are strictly larger. Orientations differing only on unencountered pairs tie throughout the trace.*

Proof. Two different orientations of a pair disagree on both incoming labels. Thus their first disagreement along the trace must occur on the first encounter of a pair. At that edge G uses zero and O uses one; incoming histories agree, so the outgoing difference $H_O - H_G$ is one. A nonrecording edge preserves this difference. A recording edge updates a positive difference d to

$$2d + b_O - b_G \geq 2d - 1 \geq 1.$$

This includes later visits to either edge of the same pair. The history gap therefore never closes. Before the first disagreement, primitive operations and costs agree, regardless of their private-control names. At that disagreement, Lemma 46.2 gives a strictly later alternative arrival. Afterward, Lemma 46.1 makes every alternative recorded macro strictly longer, and prime-cost monotonicity makes every nonrecording edge no faster. The cumulative strict gap persists. If there is no encountered disagreement, all histories and clocks coincide. The termination ranks already proved ensure that every finite arrival clock is defined. $\square$

Equivalently, after k recording events,

$$H_k = 2^k h_0 + \sum_{j=1}^k b_j 2^{k-j}.$$

The greedy feasible bit word is lexicographically smallest. Its leading favorable difference outweighs all later forced opposite differences, including repeated pair visits. This explains the history optimum; the operation-embedding lemma is still essential for the literal-clock optimum.

Corollary 47.2 (Every clean startup is optimal within this class). *For every original L, R, T, C loader input, the present source minimizes history and literal arrival time at every normalized startup prefix among all 2^{233} orientations of the fixed templates. Exactly 2^{227} static orientations attain the final startup minimum, counted as label assignments rather than machines modulo isomorphism.*

Proof. The normalized startup word for arbitrary natural T is

entry; (v0000p,v0000d)^T; v0000z; n0001e0; n0001e1; n0001e2; n0001e3.

The entry first encounters its pair, so it gets zero; later clearing decrements use its companion and get one. The positive-test edges `v0000p` are nonrecording. The next five distinct first encounters are the zero exit and the four merge edges. Thus the same six zero choices in Table 6 are greedy for every T , independently of L, R, C . Theorem 47.1 applies. All six pairs are encountered even for $T = 0$, so a change on any of them is strictly worse by the final boundary. The other 227 pairs are unused during startup and can be oriented freely. $\square$

Consequently (45.1) is the exact minimum startup clock in this class for every $T = 0$ input; 138 is its attained empty-input minimum. This does not claim that the present orientations of the other 227 pairs are optimal for later TM prefixes or simultaneously for all full inputs. The theorem about this particular source and its predecessor remains Theorem 42.1.

Given a finite trace, the greedy assignment is computed by one scan with a set of encountered pairs. The resulting literal machine contains no simulation callback or halting oracle. Nothing here supplies a uniformly terminating procedure that compiles an optimal completed assignment for an arbitrary infinite computation: an ongoing scan need not certify which unseen pairs will never occur. No global machine/runtime optimality, raw-CA optimum, or optimality beyond the fixed templates is asserted. The separate proof and independent audit are bundled in `reproducibility/orientation-addendum/`; they do not change any frozen source or recorded execution bytes.

48 The genuine equal history slowdown

The four companion edges in Table 6 receive one instead of zero. From *equal* starting history h , each such macro now costs two extra five-counter steps and ends at history $2h + 1$ instead of $2h$. The prime encoding can greatly magnify this local difference.

For example, execute the `v0303z` macro from `tm_I0_0_restore` with five-counter input $(0, 0, 0, 0, 0)$ and $C = 1$. The physical starting pair is $(1, 0)$ in each machine. Direct literal replay gives

Machine	Literal steps	Final history	Final physical pair
Frozen predecessor	28	0	$(1, 0)$
New source	104	1	$(7, 0)$

Both reach `tm_A0_pop0`; both paths were traversed. This rules out universal pointwise speedup from arbitrary internal states with equal histories. It does not contradict Theorem 42.1: corresponding clean-loader runs arrive here with old history already at least 15 larger.

More generally, along any fixed normalized path with k recording events, its history difference from an equal initial history is

$$H_{\text{new}} - H_{\text{old}} = \sum_{j=1}^k (b_{n,j} - b_{o,j}) 2^{k-j}. \quad (48.1)$$

This expression may have either sign. The clean-loader theorem uses the specific favorable startup and its persistent positive old-minus-new gap; it does not infer a time comparison for arbitrary words from their final history values alone.

49 The executed 348 step prefix

The literal interpreter evaluates the actual saved branch guards, selects the unique enabled branch, changes one physical counter by its serialized delta, and advances one control per

step. It never jumps over a macro. It traversed empty startup and one complete following TM transition in one continuous run. All 348 source rows are recorded in the bundled `empty-through-first-tm-literal-trace.json`.




Figure 6: The actual continuous source execution. Labels below the arrows refer to a different simulation layer, not alternative counts of physical steps. Exactly one TM transition occurs after the first cut.

Arrival boundary	Five steps	Literal increment	Literal total	(A, B)
<code>init_clear_T</code>	4	22	22	(1, 0)
<code>n0001_0</code>	8	28	50	(1, 0)
<code>n0001_1</code>	12	22	72	(1, 0)
<code>n0001_2</code>	16	22	94	(1, 0)
<code>n0001_3</code>	20	22	116	(1, 0)
<code>tm_A0_pop0</code>	24	22	138	(1, 0)
<code>tm_B0_pop0</code>	46	210	348	(7, 0)

Table 9: A compact index into the saved trace. “Five steps” is cumulative at the five-counter layer. “Literal increment” is the duration since the preceding listed boundary. No individual physical row is omitted from the machine-readable trace.

At the empty input, each identity costs one literal step and every zero test in the startup costs seven. The first history macro has one identity and three tests, costing 22; the second has four tests, costing 28; each of the last four again costs 22. Their sum is 138.

The first TM transition has four independent three-counter instructions and 22 five-counter steps. Its first four history macros have zero tags and cost 28 literal steps each. The last has bit one and its six primitive costs are 1, 7, 7, 17, 35, 31, totaling 98. Thus the literal duration is $4 \cdot 28 + 98 = 210$. It reaches state B scanning zero, with $(L, R, T, H, W) = (0, 0, 0, 1, 0)$, hence physical (7, 0). This is a finite prefix, not a completed universal program or evidence of efficient general simulation.

Executed segment	Zero updates	Decrements	Increments
Startup	100	19	19
First TM transition	134	35	41
Continuous prefix	234	54	60

Table 10: Serialized branch census. The rows sum to 138, 210, and 348, respectively. “Zero update” includes tests and identities.

The producer also traversed five other clean startups, at $A = 2, 3, 13, 72, 78$, with respective literal clocks 214, 290, 1130, 6118, 6650. The independent checker traversed 13 clean startups at $A = 1, 2, 3, 4, 6, 9, 13, 17, 26, 27, 81, 243, 351$; for $A = 351$ the exact count is 29,706. These finite cases corroborate Proposition 45.1; the 24-operation decomposition proves it for every allowed A .

50 The binary compiler and the limits of the clock claim

The unchanged strict compiler uses the numerical parameters

$$\begin{aligned} m &= 122622, & p &= 66066, & a &= 75495, & J &= 0, \\ D &= 509508, & S &= 1019018, & Z &= 20380380. \end{aligned}$$

Its indexed construction still has 269,291,358,255 local involution factors and radius upper bound 3,292,955,588,459,274,804. The binary alphabet and five-particle encoding are unchanged. Empty input is loaded at positions

$$\{-20380381, 0, 1019018, 1019019, 20380380\}.$$

The source table differs, so this is a new indexed CA rule despite identical geometry. No equality or conjugacy of its global map with the predecessor’s map is asserted.

For the unchanged compiler, a source zero update takes one CA microedge. A source move with pre-move selected physical counter c and $\delta \in \{-1, 1\}$ takes

$$3 + 2(Z + c) + \delta - 4S \tag{50.1}$$

microedges. Summing this theorem-derived expression over the actually executed source trace gives:

Source segment	Predicted CA microedges
Empty startup	1,394,018,396
First TM transition only	2,788,036,936
Through first TM transition	4,182,055,332

These are exact predicted CA clocks. No CA microedge run was traversed and no full factor array was allocated. In particular, a count of fewer literal steps does not by itself compare sums of (50.1), whose costs also depend on intermediate physical counters. This report proves no old/new raw-CA speed comparison and gives no uniform speedup ratio.

The source properties needed for the earlier binary composition are retained: partial injection, predecessor-free start, terminal halt, and nonblocking infinite continuation on clean inputs. Thus the existing periodicity argument applies separately to the new CA instance. If a finite source run has total CA microedge clock Θ , its signed orbit has least period $2\Theta + 2$; if the source continues forever, the initial encoded configuration never returns. The no-incoming start and partial injection exclude a repeated unsigned forward node. This is proof inheritance through a newly verified source, not an execution of a universal CA instance. The separate four-mass lower-bound theorem of Report 12 is not re-proved or needed for the startup comparison here.

51 Reproducibility and provenance

The portable bundle has three roles: this editable L^AT_EX report and its PDF; the exact source, proofs, traces, and reconstruction checks in `reproducibility/`; and release-wide byte manifests. All runtime references are relative to the extracted bundle. Verification requires Python 3 and its standard library, and no network. Rebuilding this PDF additionally requires a L^AT_EX installation with the packages named in its preamble.

The pinned source dependency is the 528-instruction three-counter program for the Neary–Woods 15-state, two-symbol universal TM [1]. Its source rows, normalization, and input encoding

are unchanged. The published finite-input universality theorem is an inherited mathematical dependency. The bundled transition table and entry-by-entry comparison do not purport to re-prove that published theorem. The earlier primary visual check covered printed pages 112 and 121, including all 30 transition cells and the initial scanned blank. The undefined transition is state J scanning one. No new third-party PDF is bundled.

The history and prime constructions follow Morita’s reversible-counter method [2], but their exact local graphs are independently reconstructed and proved here. The prior primary scan was not visually verified: the available endpoint returned HTML and screenshots failed. No claim relies on resolving an ambiguous OCR symbol. These inherited provenance limits remain explicit.

The checks have distinct logical jobs:

1. Reconstruct the primitive and normalized three-counter graph from pinned input; verify the complete source simulation and its affine paths.
2. Verify every per-pair label bijection, reconstruct the full five-counter graph, and reconstruct every expanded two-counter prime graph.
3. Check complete edge multisets, private-state separation, signed serialization, global domain/image separation, and START/HALT conditions.
4. Independently derive the exact four relabelings from actual normalized rows and compare against mandatory pinned predecessor inputs.
5. Traverse the startup and 348-step source prefix; check the nonzero- T five-counter formulas, equal-history counterexample, and 336 suffix embeddings.
6. Rebuild all seven generated source outputs byte-for-byte, then replay all checks both normally and with Python’s optimization flag enabled.

The checker uses explicit failure conditions rather than relying on assertions that disappear under `-O`. The broader macro audit includes 2,330 history and 27,743 prime sample runs; the targeted relabel audit adds 1,398 history traces and eight five-counter prologues for $T = 0, \dots, 7$. Finite sample counts are supplemental evidence. The unbounded proofs are the checked complete graphs together with their invariants and ranks.

Consult `reproducibility/README.md` for the exact replay commands and `PORTABILITY.md` there for the adaptation ledger. Baseline files are bundled and SHA-pinned; a missing baseline is a failure, not a reason to skip the comparison. Any retained historical receipt is explicitly bounded to its original role and labeled if adapted. Fresh portable replay receipts attest to the shipped checkers, not to an undocumented external workspace.

The orientation addendum is a separate independently checked mathematical argument. Its two notes and seal are verified as file dependencies; this hash check does not mechanically prove their theorems or add literal-execution coverage.

The complete trace SHA-256 is

```
6f6374db5fb0f8ce1de3afb628cbee1e46c354ad98d88b0d2c278601f6005e4d
```

The new reversible-five graph SHA-256 is

```
2f426dda9743a80d1677bf367f4c75b1522efbd31cd6f07ab59ff7643b4ed31f
```

The outer manifest covers the actual shipped files and excludes only itself and its detached checksum to avoid self-reference. The archive has its own separately recorded SHA-256. No developer logs, caches, absolute workstation paths, or optional cross-release reads are part of the portable package.

52 What is resolved and what remains useful

Does choosing zero lose reversibility? No. The colliding pair still has distinct zero/one labels, and every global join is rechecked. A repeated zero history value at different controls is not a collision of complete configurations.

Does this quietly narrow the loader? No. Equations (45.2)–(45.3) cover all initial natural scratch, and all original L, R, C choices survive. The small formula (45.1) is expressly restricted to $T = 0$.

Is the improvement only a startup benchmark? No. The history-gap recurrence and chronological embedding prove a whole-prefix comparison at corresponding boundaries from every clean loader input. However, later history can still become enormous, so the theorem supplies no practical runtime guarantee for arbitrary universal programs.

Can every internal state be made faster? This particular change cannot: the fully executed companion example costs 104 instead of 28 steps at equal history. For a specified finite workload trace, Theorem 47.1 gives the optimal orientation within these fixed templates. That conclusion does not optimize arbitrary internal inputs or changes to the recorder itself.

What would make the next computation useful? A sensible next target is another explicitly bounded literal TM prefix with complete trace accounting. One should first compute its five-counter and prime-expanded clock to determine whether literal replay is feasible. This report stops at the one transition that was actually executed; no pending run is presented as a result.

What remains at the CA layer? A lazy indexed compiler or evaluator would avoid eagerly allocating hundreds of billions of factors, but would need its own semantic and indexing audit. A genuine old/new CA clock comparison needs a separate argument controlling intermediate physical counter values and move counts, not just Theorem 42.1. Those are concrete next questions, not claimed consequences of the present proof.

[write] Both questions are answered in this report: Report 18 (second half of this Part) gives the old/new comparison in the automaton’s clock (Theorem 53.2), and Report 19 (Part IV) the lazy evaluator (Theorem 65.1).

Bottom line. A local, admissible choice of four history labels preserves the complete original semantics and geometry while removing a specific catastrophic clean-startup cost. The executed result is a 348-step literal prefix. The stronger unbounded result is paired clean-input boundary domination. Their evidence and limitations are deliberately kept distinct.

Cellular clocks for a literal reversible source

Exact macro costs and first encounter domination

Report 18 Separate proof and reproducible audit

3 October 2026

Abstract of Report 18, as delivered. The four history-tag swaps in Report 17 improve not only literal two-counter arrival times but also the cellular-automaton clock of its fixed strict binary compiler. The key identity is $\Theta = \lambda M + \text{TV}(A^2) + \text{TV}(B^2) + U$, with $\lambda = 36,684,691$, moving-row count M , and zero-update count U . Exact whole-prime macro costs are increasing in their enabled encoded input. A proof that all 2^{233} static history orientations have the same compiler constant then lifts the recorder's operation embedding to a CA-clock comparison. For a fixed finite normalized trace, assigning zero to each pair's first encountered edge minimizes every boundary arrival time at once. Exactly 2^{227} assignments attain the complete clean-startup minimum for every promised loader input. For zero initial scratch, that minimum has an explicit quadratic and floor formula. The empty startup takes 138 executed literal rows and exactly 1,394,018,396 theorem-derived CA microedges. The continuous executed source prefix has 348 rows and covers one TM transition; no CA orbit or giant rule array was executed. An independent graph reconstruction and normal and optimized Python replays support the proof. The portable bundle preserves the exact scientific source bytes and includes a mandatory pinned predecessor reconstruction.

What changes. This report settles the CA-clock comparison left open by Report 17 within its fixed static-orientation, recorder, prime-template, and compiler class. It does not change the Report 17 source. The two compared CA rules can differ despite equal numerical geometry.

Evidence and scope. These are mathematical proofs with independent executable audits, not proof-assistant formalizations. No novelty, optimal universal-machine size, unrestricted speedup, or host-runtime claim is made. A clock unit is one iteration of the compiled CA on its admissible forward plus orbit, not one evaluation of an individual factor of its local rule.

[write] Delivered as `Cellular_Clock_Domination_Package.zip` (batch 82, archive 04; arrival `db37d18c8`), dated 3 October 2026; files placed by `bca6383e9` with the prefix given in Appendix A. Report 18 continues Report 17 (first half of this Part) and settles the cellular-clock comparison that Report 17 left open.

53 The comparison and its domain

The normalized three-counter program has data $(L, R, T) \in \mathbb{N}^3$. Its fixed recorder adds history H and work W . A five-counter boundary is represented in the literal two-counter machine by

$$(A, B) = (C2^L 3^R 5^T 7^H 11^W, 0), \quad C \geq 1, \quad \gcd(C, 2310) = 1. \quad (53.1)$$

The clean-loader domain is $H = W = 0$; it includes every natural initial scratch value T , not just $T = 0$. Every literal row is an increment, enabled decrement, zero test, positive test, or identity. Tests cost time.

Let $\mathcal{O}$ consist of the 2^{233} static orientations of the 233 incoming collision pairs. An orientation gives distinct bits zero and one to the two members of each pair. The normalized program, recorder, prime order $(2, 3, 5, 7, 11)$, literal macro templates, and strict binary-five compiler are fixed. Every member retains the recorder's global partial injection and its enabled-macro completion property. Arrival times below are at completed normalized-source edges, not at arbitrary internal rows.

Theorem 53.1 (First encounter prefix optimality). *Fix a finite enabled normalized trace, common initial data, common history $h_0 \in \mathbb{N}$, clean work $W = 0$, and common cofactor C as in (53.1). Choose $G \in \mathcal{O}$ to give bit zero to the first encountered edge of each used pair. Then G minimizes cumulative CA arrival time at every normalized boundary simultaneously. A competitor ties through a prefix exactly when it agrees with G on all pairs encountered in that prefix. At its first encountered disagreement it becomes strictly slower, and it remains strictly slower at every subsequent boundary.*

[write] The literal-clock statement is Theorem 47.1 of Report 17.

Theorem 53.2 (The frozen optimized and predecessor sources). *On every common input $(\text{START}, C2^L3^R5^T, 0)$ in the clean-loader domain, the Report 17 source reaches every corresponding normalized boundary no later in CA microedges than the Report 16 predecessor. It arrives strictly earlier by the first TM cut, and at every later corresponding boundary. The assertion holds for every finite prefix, whether or not the run halts.*

[write] The literal-clock statement is Theorem 42.1 of Report 17; this theorem answers the comparison that Report 17 left open (Section 52).

A TM cut requires the designated TM control and $T = 0$. A loop visit to a cut label with nonzero scratch is not such a cut. Both sources have the same normalized data trace, including the same TM transitions.

These theorems compare different fixed rules at matching semantic boundaries. They do not order arbitrary equal-history internal states or all intermediate physical configurations. They assert no equal-time alignment, conjugacy, fixed speedup ratio, or globally optimal construction. The finite trace in Theorem 53.1 is supplied; no halting decision procedure or uniform discovery of an arbitrary infinite run's first encounters follows.

54 An exact total variation identity

The inherited strict compiler subdivides a moving literal row, with old selected counter c and update $\delta \in \{-1, +1\}$, into dispatch, outward travel, endpoint move, return travel, and commit. Their durations are

$$1, \quad Z + c - 2S, \quad 1, \quad Z + c + \delta - 2S, \quad 1.$$

Thus its exact microedge clock is

$$\tau(c, \delta) = 3 + 2(Z + c) + \delta - 4S = \lambda + \left| (c + \delta)^2 - c^2 \right|, \quad \lambda = 3 + 2Z - 4S. \quad (54.1)$$

For an increment $c \geq 0$, the absolute difference is $2c + 1$. For an enabled decrement $c \geq 1$, it is $2c - 1$. A disabled decrement is not covered. Every zero-update row, including each successful test and identity, takes exactly one microedge.

[write] Equation (54.1) is Report 15's clock (5) (Section 8) and Report 16's (29), written as a total variation.

Proposition 54.1 (Exact path clock). *For every finite enabled literal path, let M count moving rows and U count zero-update rows. If (A_t, B_t) are its actual physical counters, then*

$$\Theta = \lambda M + \sum_t |A_{t+1}^2 - A_t^2| + \sum_t |B_{t+1}^2 - B_t^2| + U = \lambda M + \text{TV}(A^2) + \text{TV}(B^2) + U. \quad (54.2)$$

Proof. Sum (54.1) over moving rows and add one for each zero update. Only the selected physical counter changes in a moving row, while neither changes in a test or identity. Every row is counted once. $\square$

Total variation must not be replaced by the endpoint difference of squares. For example, a counter moving $N \rightarrow 0 \rightarrow N$ contributes $2N^2$, although its endpoints agree. This distinction is why literal instruction-count improvement alone was insufficient to establish CA-clock improvement.




Figure 7: A return to the same counter value still contributes positive clock.

Writing m for controls and p_{move} for moving rows, the compiler has

$$D = 2(m + 2p_{\text{move}}), \quad S = 2D + 2, \quad Z = 40D + 60 + 2J, \quad \lambda = 72D + 115 + 4J > 0.$$

For the present class, $m = 122622$, $p_{\text{move}} = 66066$, $J = 0$, $D = 509508$, $S = 1019018$, and $Z = 20380380$, so $\lambda = 36,684,691$. Section 56 proves that these values are common to every compared orientation.

55 Whole prime macros from their literal paths

A five-counter operation at prime $p \in \{2, 3, 5, 7, 11\}$ begins with physical counters $(N, 0)$, $N > 0$. Other encoded exponents are arbitrary naturals. An enabled decrement or positive test requires $p \mid N$; an enabled zero test requires $p \nmid N$. We count the entire macro through its restored five-counter boundary, including its destination restorer.

55.1 Increment and enabled decrement

For increment, the first phase transfers N units from A to B ; the second removes each B unit and increments A exactly p times. Hence $A : N \rightarrow 0 \rightarrow pN$ and $B : 0 \rightarrow N \rightarrow 0$. The exact counts are

$$M = (p + 3)N, \quad \text{TV}(A^2) + \text{TV}(B^2) = (p^2 + 3)N^2, \quad U = 4N + 3.$$

The three isolated tests are entry $B = 0$, transfer exit $A = 0$, and final $B = 0$. The transfer and return iterations each contribute two positive tests. Therefore

$$I_p(N) = (p^2 + 3)N^2 + \lambda(p + 3)N + 4N + 3. \quad (55.1)$$

For enabled decrement put $N = pQ$. The return loop removes p units of B per iteration and adds one to A , so $A : N \rightarrow 0 \rightarrow Q$ and $B : 0 \rightarrow N \rightarrow 0$. Every individual decrement in a p -block is charged. Divisibility prevents an incomplete block. We obtain

$$\begin{aligned} M &= 3N + Q, & \text{TV}(A^2) + \text{TV}(B^2) &= 3N^2 + Q^2, & U &= 2N + 2Q + 3, \\ D_p(N) &= 3N^2 + Q^2 + \lambda(3N + Q) + 2N + 2Q + 3. \end{aligned} \quad (55.2)$$

55.2 Successful zero and positive tests

Write $N = pQ + r$, $0 \leq r < p$. The scanner decrements A unit by unit, tracks residue in finite control, and increments B once per full block. It exits at $(0, Q)$. Restoration first adds back r , then adds p units for each removed unit of B , ending at $(N, 0)$. Consequently

$$\begin{aligned} M &= 2(N + Q), & \text{TV}(A^2) + \text{TV}(B^2) &= 2(N^2 + Q^2), & U &= 2(N + Q) + 3, \\ T_p(N) &= 2(N^2 + Q^2) + 2(\lambda + 1)(N + Q) + 3, & Q &= \lfloor N/p \rfloor. \end{aligned} \quad (55.3)$$

The zero updates are one entry test, $N + Q$ scanner positive tests, one scanner exit, $N + Q$ restoration positive tests, and one final test. The formula includes restoration exactly once and is valid when $Q = 0$. It applies only to the selected successful branch. A wrong singleton test blocks; we assign it no successful-boundary clock. An identity costs one.

These formulas evaluate exact paths in the serialized source. No repeated arithmetic, division, residue test, or restoration becomes a unit-cost operation of the simulated machine.

56 Monotonicity and the common compiler constant

Lemma 56.1 (Whole macro monotonicity). *For every real $\lambda > 0$, a fixed prime and a fixed successful primitive, the whole macro clock is nondecreasing in its positive encoded input, restricted to that primitive’s enabled domain. It is strictly increasing for increments, decrements, and tests; identity is constant. All costs are positive.*

Proof. All nonconstant coefficients of (55.1) are positive. Substituting $N = pQ$ in (55.2) gives $(3p^2 + 1)Q^2 + [\lambda(3p + 1) + 2p + 2]Q + 3$, strictly increasing for positive Q . In (55.3), both N and $\lfloor N/p \rfloor$ are nondecreasing and the N^2 contribution strictly increases. Restriction to either enabled test domain preserves the inequality. Identity has clock one. $\square$

Comparing formulas with a common λ requires more than the observation that the five-counter row count is unchanged. In particular, the literal compiler shares restore blocks by destination. We prove the stronger invariant needed to count those blocks.

At a collision pair, flipping orientation exchanges the two original source-operation payloads feeding private bit-zero and bit-one entry controls. The two transfer graphs and their shared doubling graph remain fixed. Each original operation still appears once at its original source and has the same selected counter and symbol. Therefore the multiset of outgoing source groups, classified by prime and operation kind, is unchanged.

Every swapped source row has its own private entry destination. Exchanging the two rows merely permutes whether each entry needs restoration, and for which prime. It preserves the multiset of restorer primes. Other recorder destinations and nonrecording destinations are fixed. The argument holds after any set of other flips, so it covers all 2^{233} orientations.

Template at prime p	New controls		Literal rows	Moving rows
Identity	0		1	0
Increment or decrement	$p + 7$		$p + 10$	$p + 3$
Test group E	$2p + 2$	$2p + 3 + (p - 1)\mathbf{1}_{Z \in E} + \mathbf{1}_{P \in E}$	$2p + 3$	$p + 1$
Destination restorer	$2p + 2$		$2p + 3$	$p + 1$

Distinct prefixes keep private controls disjoint. Include the 4,520 five-counter controls once. Summing these invariant signatures gives 122,622 literal controls and 141,561 rows: 33,436 increments, 32,630 decrements, 23,429 zero tests, 52,036 positive tests, and 30 identities. There are 2,330 shared restorers, with counts (116, 116, 234, 932, 932) at primes (2, 3, 5, 7, 11). The guard grammar always has class cut $J = 0$. Thus m, p_{move}, a, J , where $a = 75,495$ counts zero-update rows, and hence D, S, Z, λ , are identical.

The indexed rules and row order may differ. The common symbolic factor count and radius bound do not identify the rules or constitute an executed CA construction. All-orientation invariance follows from the local proof; 233 individual flips and the all-flipped check corroborate it.

57 The recorder embedding charges complete CA macros

A recording edge with incoming history h and assigned bit b has word

$$\begin{aligned} &\text{source operation; } W = 0; \\ &(H > 0, H-, W+, W > 0)^h; \quad H = 0; \\ &(H+, H > 0)^b; \quad (W > 0, W-, H+, H > 0, H+, H > 0)^h; \quad W = 0. \end{aligned}$$

Here every displayed test is a separate five-counter primitive. The recorder ends at $(H, W) = (2h + b, 0)$. Compare a higher run with history $h + d$ and bit a to a lower run with history h and bit b , with equal data and $d \geq 1$.

Lemma 57.1 (Strict recording clock domination). *For all four bit pairs (a, b) , the higher-history recording edge has a strictly larger whole CA clock. At equal history, bit one costs strictly more than bit zero, and equal bits have equal clocks.*

Proof. Construct an order-preserving injection of lower operations into higher operations, matching the same symbol and selected prime:

1. Match the source operation, entry test $W = 0$, and the first h transfer iterations. The higher encoded input is 7^d times the lower.
2. Skip the higher run's d surplus transfer iterations and match $H = 0$. The encoded ratio is now 11^d .
3. If $a = b = 1$, match preparation at ratio 11^d . If $(a, b) = (1, 0)$, skip the higher preparation. If $(a, b) = (0, 0)$, neither has preparation. In these cases skip d higher doubling iterations.
4. In the remaining case $(a, b) = (0, 1)$, skip $W > 0, W-$ in the first higher doubling iteration. Match the lower preparation $H+, H > 0$ to that iteration's first such pair, at ratio $11^{d-1} \geq 1$. Finish that higher iteration and the following $d - 1$ surplus iterations.
5. Both runs now have work $W = h$. The history gap is $2d + a - b \geq 1$. Match their remaining h doubling iterations and final test at encoded ratio $7^{2d+a-b} \geq 7$.

Every lower operation is matched exactly once, both matched operations are enabled, and the higher encoded input is no smaller. Lemma 56.1 compares their complete literal expansions in CA time, using the common λ proved above. Unmatched higher operations have positive costs; in particular its surplus transfer iterations exist. This makes the total inequality strict. At equal history and bits, the words coincide. At equal history with bits one and zero, skip the bit-one preparation and match the doubling suffix and final test at history gap one, again strictly. $\square$

[write] Lemma 46.1 of Report 17 is the same embedding for the literal clock; the startup history used below is Report 17's equation (45.2), not re-derived.

For a nonrecording edge the same data operation at larger history has nondecreasing cost; identity may tie. The proof aligns whole five-counter primitives. It does not require pointwise alignment of the intermediate physical counters inside their literal prime expansions.

58 Prefix optimality and the four startup swaps

Proof of Theorem 53.1. Before the first encountered differing pair, the two histories, operations, and clocks coincide. The first differing edge must be that pair's first encounter: swapping orientation changes both of its edge labels. The first-encounter orientation uses zero there, while its competitor uses one. Lemma 57.1 creates history gap one and a strict clock delay.

At each later recording edge the competitor's history gap satisfies

$$d' = 2d + b_{\text{competitor}} - b_G \geq 2d - 1 \geq d \geq 1. \quad (58.1)$$

The higher-history macro costs strictly more, for either pair of current bits. A nonrecording edge preserves the history gap and has no smaller clock. Therefore the cumulative delay cannot shrink. This proves every prefix inequality and the stated equality characterization, including repeated visits. Every finite enabled macro completes, so all compared finite arrival times are well-defined. $\square$

If a finite trace encounters k distinct collision pairs, precisely 2^{233-k} orientations minimize its complete-trace clock; all of them also minimize every prefix clock. Unencountered assignments cannot affect the trace. This counts static assignments, not isomorphism classes of machines.

For every clean prologue, the first-encounter members are the same six: `entry`, `v0000z`, `n0001e0`, `n0001e1`, `n0001e2`, and `n0001e3`. The Report 17 orientation assigns zero to all six. The first two already had zero in its predecessor; exactly four pairs were swapped. The other 227 assignments are irrelevant to startup. Consequently exactly 2^{227} orientations attain the complete startup minimum for every clean-loader input simultaneously.

Proof of Theorem 53.2. The two frozen sources follow the same normalized data/control trace. Their histories agree until the first changed startup assignment, which has old bit one and new bit zero. Its old clock is strictly greater. The common clearing loop ends with history $2^T - 1$. The following five recording edges give

$$H_{\text{old}} = 32(2^T - 1) + 15, \quad H_{\text{new}} = 32(2^T - 1). \quad (58.2)$$

Thus by the first TM cut the old history exceeds the new by exactly 15. Thereafter (58.1), now with the old and new bits, preserves a gap of at least 15. Recording macros cost more in the old run and nonrecording macros no less. Induction gives the claim for every later finite boundary. $\square$

This second theorem does not assume that the Report 17 orientation is first-encounter optimal on every later computation. Its all-later-boundary comparison follows from the already established positive history gap. The scope remains paired clean inputs; unrelated equal-history internal states are not a substitute for that hypothesis.

59 The exact minimum when initial scratch is zero

For $T = 0$ write $A = C2^L3^R$ and $q_p = \lfloor A/p \rfloor$. The optimized five-counter startup has five identities and nineteen successful zero tests: one at prime 5, six at prime 7, and twelve at prime 11. All occur at the same encoded integer A because all six first-encounter bits are zero. The loader's coprimality and $T = H = W = 0$ make these tests enabled.

Proposition 59.1 (Closed clean startup clock). *The exact minimum CA startup time in $\mathcal{O}$ for $T = 0$ is*

$$\begin{aligned} \Theta_0(A) = & 38A^2 + 2q_5^2 + 12q_7^2 + 24q_{11}^2 \\ & + 2(\lambda + 1)(19A + q_5 + 6q_7 + 12q_{11}) + 62. \end{aligned} \quad (59.1)$$

Proof. The startup costs $5 + T_5(A) + 6T_7(A) + 12T_{11}(A)$. Substitute (55.3) and collect terms. The constant is $5 + 3(1 + 6 + 12) = 62$. First-encounter optimality makes this achieved value the class minimum. $\square$

At empty input, $A = 1$, all three quotients vanish, so

$$\Theta_0(1) = 38\lambda + 138 = 1,394,018,396. \quad (59.2)$$

The source actually takes 138 literal rows: 38 moves and 100 zero updates. Their total square variation is 38, agreeing with $\lambda \cdot 38 + 38 + 100$. The large CA count is derived from the compiler travel-time theorem. The CA itself was not iterated that many times.

Executed source segment	Literal rows	Derived CA microedges
Empty startup	138	1,394,018,396
Next TM transition	210	2,788,036,936
Continuous prefix	348	4,182,055,332

The third row is the continuous union of the first two segments, not a third disjoint execution. The first transition ends at control `tm_B0_pop0` with physical counters (7, 0), representing one unit of history and zero data. The earlier cut is `tm_A0_pop0` at (1, 0). Only one TM transition beyond startup is certified by this saved trace.

Formula (59.1) must not be used for $T > 0$. Although the same six choices minimize all promised startups, positive scratch drives history growth according to (58.2). Optimality within a fixed class does not make those clocks small.

Input cost boundary. Arithmetic performed by the simulated machine is fully charged. Constructing the initial integer $C2^L3^R5^T$ and writing its spatial encoding are external input preparation. The host arithmetic used to evaluate a formula is not claimed to be CA computation.

60 An exact formula for the recorded suffix

The following formula is useful for exact arithmetic without enumerating an enormous literal or CA path. It excludes the original normalized data operation. Let $K = C2^L3^R5^T$ be the data factor *after* that operation; then $\gcd(K, 77) = 1$. The suffix starts at $(H, W) = (h, 0)$ and uses bit b . Here K is a data factor, not an instruction horizon.

For one transfer iteration, let X be the encoded integer immediately after its $H-$ operation. For one doubling iteration let Y be the encoded integer immediately after $W-$. If V is the input to bit-one preparation, the whole phase costs are

$$\mathcal{T}(X) = 616X^2 + (76\lambda + 60)X + 12, \quad (60.1)$$

$$\mathcal{D}(Y) = 8208Y^2 + (266\lambda + 208)Y + 18, \quad (60.2)$$

$$\mathcal{P}(V) = 152V^2 + (26\lambda + 20)V + 6. \quad (60.3)$$

Each is obtained by summing (55.1)–(55.3) for every primitive in that phase. For example, preparation is $I_7(V) + T_7(7V)$; its coefficients are $52 + 100$, $10 + 16$, $4 + 16$, and $3 + 3$.

Define integer geometric sums, including the value zero at $h = 0$,

$$U_1 = \frac{11^h - 7^h}{4}, \quad U_2 = \frac{121^h - 49^h}{72}, \quad V_1 = \frac{49^h - 11^h}{38}, \quad V_2 = \frac{2401^h - 121^h}{2280}.$$

The exact recorded-suffix clock is

$$\begin{aligned} \mathcal{R}(K, h, b) = & T_{11}(K7^h) + 616K^2U_2 + (76\lambda + 60)KU_1 + 12h \\ & + T_7(K11^h) + b[152K^2121^h + (26\lambda + 20)K11^h + 6] \\ & + 8208K^249^hV_2 + (266\lambda + 208)K7^bV_1 + 18h \\ & + T_{11}(K7^{2h+b}). \end{aligned} \quad (60.4)$$

To see the sums directly, the transfer after-decrement inputs are $X_i = K7^{h-i-1}11^i$ for $0 \leq i < h$. Their sums are KU_1 and K^2U_2 after squaring. The doubling after-decrement inputs are $Y_j = K7^{b+2j}11^{h-j-1}$ for $0 \leq j < h$, giving $K7^bV_1$ and $K^249^bV_2$. Preparation has $V = K11^h$. Add the initial $W = 0$, intermediate $H = 0$, and final $W = 0$ tests. This proves (60.4), including the empty iteration blocks.

The producer compares this formula to explicit primitive words in 156 cases. The independent auditor separately derives the three coefficient vectors by exact rational arithmetic. Neither calculation substitutes finite tests for the all-natural derivation or adds a unit-cost arithmetic operation to the simulated machine.

61 What was executed and what was inferred

The saved continuous run traverses actual branches of the serialized literal source. Both producer and independent checker replay every control, counter pair, branch name, and per-row clock. A separate sum over the saved five-counter traces gives the same CA totals.

Segment	Moves	Zero updates	$TV(A^2)$	$TV(B^2)$	Five primitives
Startup	38	100	38	0	24
Next TM transition	76	134	282	4	22
Continuous prefix	114	234	320	4	46

For example, the whole prefix satisfies

$$114 \cdot 36,684,691 + 320 + 4 + 234 = 4,182,055,332.$$

The step totals are actual literal-row executions. The microedge totals are exact theorem-derived predictions. Neither the giant local-factor array nor a CA trajectory was allocated or executed. The independent macro tests also traverse literal rows, but these are separate bounded validation cases, not additional progress in the universal computation.

A separate old/new startup check executes 142 predecessor and 24 optimized five-counter rows. Summing whole-prime clocks gives respectively 126,594,455,831,808,973,674,445,902 and 1,394,018,396 CA microedges. The predecessor's full literal startup and both CA orbits remain unexecuted; the audit independently checks this arithmetic and preserves the complete five-counter traces.

The independent audit reconstructs all 5,451 five-counter rows and all 141,561 literal rows, checking complete edge multisets and serialized branches. Its graph reconstruction uses neither producer graph routines nor the compiler. The actual rows are reconstructed from the preceding graph layer and certificate private names, with certificate endpoints, operations, primes, and coverage checked against that layer.

Its main finite test battery includes:

- 21,277 successful and 10,136 blocked macro traces, traversing 2,795,605 actual literal rows. Successful traces compare moving counts, square variation, and zero-update counts separately.
- 8,112 adjacent enabled-input monotonicity comparisons, including floor transitions; 975 explicit recorder embeddings with 62,361 matched primitive operations.
- All 233 individual pair flips, the all-flipped orientation, and the actual predecessor orientation, with the source-group/restorer census.
- 4,096 startup comparisons from 64 inputs and all 64 startup assignments, checked at 36,864 normalized boundaries, including the exact equality and strictness conditions.

The startup grid has $L, R \in \{0, 1\}$, $T \in \{0, 1, 2, 3\}$, $C \in \{1, 13\}$, and $h_0 \in \{0, 1\}$. The embedding grid uses data factors 1, 13, 30, histories 0 through 12, positive gaps 1 through 6, all four bit pairs, and equal-history comparisons. These tests do not enumerate 2^{233} assignments. Structural invariance and the embedding argument supply the unbounded proof. Every check uses explicit exceptions, so Python optimization does not disable it.

62 Connection to finite horizon timing certificates

The inherited direct-primitive shared-offset certificate provides the appropriate Diophantine connection [14]. Fix a literal source instruction horizon n as compiler syntax. At each layer its natural selectors $e_{t,\alpha}$ choose one original branch, and its shifted postcounter and offset variables define actual pre-counter expressions $C_{t,j}$. Those expressions are constants at $t = 0$ and sums of two linear witness coordinates thereafter.

The certificate already retains an exact timing witness Θ and the single residual

$$\mathcal{E}_{\text{time}} = \Theta - \sum_{t=0}^{n-1} \sum_{\alpha} e_{t,\alpha} \tau_{\alpha}(C_{t,j_{\alpha}}). \quad (62.1)$$

For a zero-update branch $\tau_{\alpha} = 1$; for a moving branch it is $\lambda + 2C_{t,j_{\alpha}} + \delta_{\alpha}$. Each summand has degree at most two, so this residual has degree at most two and its square degree at most four. The exact-clock extension adds one retained natural variable and one squared residual to the core. Supplying a fixed clock externally instead adds only the residual. At horizon zero it forces $\Theta = 0$.

The inherited certificate's core enforces an enabled first-halt trace. Its natural witness fiber is empty or a singleton. The timing equation uniquely forces the clock and leaves that property intact. The separately paid nonnegative-real version additionally uses one selector-norm residual per layer; the timing result does not make those integrality-enforcing rows free. Its CA interpretation remains conditional on the strict compiler's simulation and observation theorem.

Proposition 54.1 is an identity for the value already computed by (62.1). The complete prime-macro formulas are an exact arithmetic evaluation shortcut. They need not be inserted as new floor, absolute-value, geometric-sum, or auxiliary-variable relations in the polynomial. In particular, this report does not silently treat those expressions as polynomial primitives.

The branch census and compiler constant are invariant under orientation. Thus the inherited fixed-horizon arity and written-term-count formulas based on that census remain the same at the *same literal horizon*. A faster corresponding computation can have a different literal horizon; that fact alone is not a same- n certificate-size reduction. Full coefficient encodings, square expansion, and numerical witness heights retain their existing charges. No numerical arity improvement is claimed here, and no giant universal-instance polynomial is emitted by this report.

Under the inherited signed-orbit theorem, a halting clean-loader run from the predecessor-free START, with forward first-halt clock Θ , yields least period $2\Theta + 2$; an infinite promised source run never returns to the encoded start. If the paired computation halts, Theorem 53.2 therefore gives a strictly shorter such period for the optimized instance. This is a conditional transfer through that separate theorem, not a new CA execution or an unrestricted-time unique-fiber Diophantine construction.

63 Reproducibility and scientific provenance

This release includes the editable L^AT_EX source, its PDF, and a portable Python-standard-library replay package. The scientific tables and saved traces are byte-identical to the sealed Report 17 release. The predecessor is regenerated from its bundled pinned generator and dependencies into temporary storage, then checked against mandatory expected output hashes. No external sibling release or optional fallback is needed.

From a fresh extraction, verify the complete delivered inventory:

```
python3 -B verify_release.py
```

Then run the aggregate checks:

```
cd reproducibility
python3 -B run_checks.py
```

The wrapper runs fresh normal and optimized Python copies. All required checks and input pins remain mandatory. Its temporary replay leaves the shipped files unchanged. The package documents portability adaptations separately from the original scientific bytes, including adapted script hashes and changed path/output plumbing. Individual checker commands can regenerate receipts; use the aggregate wrapper for a nonmutating replay. The proof and independent audit are ordinary mathematical documents, while their executable checks are reproducible finite verification.

The unchanged Report 17 source SHA-256 is

```
fa61d06178d718511232a80a05e02192d940fde3c273f0c202f0623253b2ade3
```

The mandatory regenerated predecessor source SHA-256 is

```
38c586706fa1069442d3e5103151adb156448b7d7c783f5c46d9d507b87fe73a
```

The independently reviewed cellular-clock proof SHA-256 is

```
f7385dc48b48fb4f3f95ca082477f36184d6b9cd83f8a97962e9edcc75f94c4c
```

The source dependency is the unchanged three-counter translation of the Neary–Woods 15-state, two-symbol universal TM [1]. Its published finite-input universality theorem is inherited, not re-proved by these bounded executions. The transition table, provenance, and all 30 entry checks are bundled. Their earlier primary-page visual inspection is historical evidence; replay checks the pinned transcription rather than reopening the paper. The counter construction follows Morita’s method [2]; the exact implemented graphs have their own reconstruction proof. The prior Morita scan was not visually verified, and no argument here relies on ambiguous OCR. No third-party full paper is shipped.

The fixed compiler proof and source schema are bundled. The retained certificate reference is proof-only: Report 18 checks its byte pin but does not rerun the separate certificate software or materialize its universal polynomial. Earlier Reports 16 and 17 remain separate, unchanged releases. The new result strengthens their stated clock knowledge without altering their historical evidence or claims.

64 Conclusions and remaining boundaries

The obstruction to a CA-clock comparison was genuine: individual literal moves have state-dependent cost, so counting fewer instructions is not enough. The total-variation identity, exact whole-macro monotonicity, and all-orientation compiler invariance supply the missing argument. The recorder embedding then proves both the first-encounter prefix optimum and the paired predecessor/optimized domination theorem.

The precise conclusions are:

1. Every matched normalized boundary in paired clean runs is reached no later by the optimized source, with strict improvement by the first TM cut and thereafter, in CA microedges.
2. First encounters determine all prefix-optimal static orientations for any given finite enabled normalized trace. The startup family fixes six choices and leaves exactly 2^{227} optimal assignments.
3. For zero initial scratch, the exact minimum is (59.1); for a general recorded suffix, (60.4) evaluates every charged phase without executing it.
4. The proof is backed by independent full-graph reconstruction, finite macro/embedding tests, and a replay of the 348-row continuous literal prefix in ordinary and optimized Python.

The result does not optimize a different recorder, prime ordering, encoding, compiler geometry, or arbitrary universal architecture. It does not show a uniform ratio between arrival times. It neither compares arbitrary internal initializations nor aligns different CA configurations at equal times. Symbolic local-factor and radius counts are not allocated objects. The universal computation represented by the saved prefix was not completed, and no CA orbit was executed.

A useful next arithmetic question is whether the full nonzero- T startup has a comparably convenient closed CA-clock expression. A different research problem is the tradeoff obtained by changing the recorder or compiler itself. Neither is needed for the theorems proved here.

Part IV

Exact lazy evaluation (Report 19)

Exact lazy finite support evaluation of the frozen reversible binary CA

Preserving an ordered product of 269291358255 factors

Report 19 Mathematical proof and portable reproducibility package

3 October 2026

Abstract of Report 19, as delivered. The frozen compiler of Report 15 defines a reversible, number-conserving binary cellular automaton by an ordered finite product of isolated local involutions. Its pinned universal source from Report 17 yields 269,291,358,255 factors, making literal factor-array construction impractical. We give a source-sized metadata representation, arithmetic random access to an individual factor, and a sparse evaluator that is exactly equal to that same ordered product on every finite binary support. Candidate completeness, exact simultaneous single-factor semantics, and rediscovery after each change establish the result even for malformed configurations. The inverse uses the exact reversed order. Costs depend on support size, source incidence, and the actual factor cascade; no constant-time or polynomial-in-log-factor-count guarantee is claimed. Differential suites, a separate audit, an explicit malformed-input cascade, and 128 literal universal startup microsteps with reverse checks accompany the proof. The steps do not complete startup or a Turing transition. The package contains the exact compiler, the exact compressed source, complete implementation, proofs, normal and optimized receipts, and an offline extracted replay. This is an implementation-equivalence result, not a new universality proof, a proof-assistant formalization, or a claim of novelty.

[write] Delivered as `Exact_Lazy_Reversible_CA_Evaluator_Package.zip` (batch 82, archive 09; arrival `db37d18c8`), dated 3 October 2026; files placed by `bca6383e9` with the prefix given in Appendix A. Report 19 evaluates Report 15’s automaton (Part I) for Report 17’s source (Part III) and answers Report 16’s second question.

65 Result and scope

A binary configuration is $x \in \{0, 1\}^{\mathbb{Z}}$. In this report its input representation is a finite support $X = \{i : x_i = 1\} \subset \mathbb{Z}$, with zero at every other site. Its numerical mass is $n = |X|$. Arbitrarily large positive and negative coordinates are permitted. A malformed input is still a legitimate finite binary configuration: no encoded-state promise, head uniqueness, separation condition, or particle-count promise is required.

Let a source accepted by the frozen compiler produce factors $g_0, \dots, g_{F-1}$ in the compiler’s exact E -then- P order. Denote the resulting map by

$$\mathcal{F}(X) = g_{F-1}(\cdots g_1(g_0(X)) \cdots).$$

Each g_i is an involution on the full binary shift, as established by the frozen isolated-swap construction [11]. The present evaluator computes the restriction of $\mathcal{F}$ to finite supports.

Theorem 65.1 (Exact finite support evaluator). *For every source accepted by the pinned reference compiler and every finite $X \subset \mathbb{Z}$, the operation `compile_lazy_source(source).step(X)` returns $\mathcal{F}(X)$. With `inverse=True`, it returns*

$$\mathcal{F}^{-1}(X) = g_0(g_1(\cdots g_{F-1}(X) \cdots)).$$

No complete factor array, travel-interval enumeration, coordinate-span enumeration, or binary local truth table is constructed. Source validation retains explicit finite class tables.

Which rule is preserved. For a fixed source the rule, all local factors, their order, and the radius bound are unchanged. Source orientation and transition relabeling are not part of this evaluator. The benchmark takes the byte-pinned Report 17 source as supplied; it does not alter that source or repeat the prior universality derivation [13]. The full-shift construction is inherited from Report 15; this implementation does not accept infinite configurations as input.

The proof below concerns all finite inputs. Tests provide executable evidence for the implementation, not a finite substitute for the theorem. Neither inverse roundtrips alone nor the universal startup sample establishes equivalence to an eager universal oracle; the latter is deliberately never built.

66 Pinned source semantics and metadata

66.1 Input schema and validation

The source schema is `reversible-two-counter-v1`. It supplies an ordered list of controls, designated start and halt controls, a nonnegative class cutoff J , and an ordered list of uniquely named branches. A branch has source and target controls, side $v \in \{-1, 1\}$, update $\delta \in \{-1, 0, 1\}$, and a finite guard expression. The side selects the updated counter; enabled transitions must keep both counters natural. The halt control has no exits. The frozen source schema and compiler proof are included in the package.

The primitive exact-type checks, guard parser, guard evaluation, and immutable branch/gate data types are reused from the pinned compiler. Its bytes are hashed *before execution*; the evaluator executes exactly the checked byte string. It does not import a reference compiler by an unpinned search-path lookup. The eager compiler is used only by tests on small sources.

Write $c = (J + 2)^2$ for the number of representative counter-class pairs. The lazy validator evaluates each branch's domain and inverse-image guard on exactly those same classes, including negative-pre-counter and negative-post-counter conditions and cutoff checks. It represents outgoing and incoming class occupancy by integer masks, indexed by control. A conflict bit is present precisely when two branches share the same domain or image class at that control.

All rows are parsed and individually validated before overlap rejection. Conflicts are then considered in the frozen control order, least class first, with domain before image at a shared class. Thus the accepted source set, branch tables, and validation error precedence agree with the reference. This ordering matters: an earlier semantic overlap cannot hide a later schema/type error that the frozen parser would report first.

66.2 Constants and stored data

Let m be the number of controls, p the number of branches with nonzero update, and a the number with zero update; put $b = p + a$. Filtering preserves original input order. Define

$$\begin{aligned} D &= 2m + 4p, & S &= 2D + 2, & B_2 &= D + 1, \\ L &= 3D + 4, & B_3 &= 4D + 5, & Z &= 10B_3 + 10 + 2J. \end{aligned}$$

A home control q (zero-based index) has signed pair gaps $2q + 1, 2q + 2$. Moving branch e has outbound gaps $2m + 4e + 1, 2m + 4e + 2$ and inbound gaps $2m + 4e + 3, 2m + 4e + 4$. In each pair, odd denotes plus and even denotes minus. These gaps uniquely identify the signed mode.

The immutable metadata contains the ordered controls and branches, filtered moving/direct branches, source/target incidence lists, class tables, arithmetic constants, counts, and a recursively immutable source snapshot. No caller-owned mutable input container is retained. The convenience encoder is

$$\text{enc}(q, c_0, c_1, \sigma) = \{-Z - c_0, 0, Z + c_1, S, S + d(q, \sigma)\}.$$

This encoder is not used to recognize arbitrary inputs during a global step.

The exact factor counts are

$$\begin{aligned} |E| &= p(4D + 15) + a, \\ |P| &= p(4D + 14) + m, \\ F &= 8pD + 29p + m + a. \end{aligned}$$

The unchanged sum-of-factor-radii bound is

$$\begin{aligned} R &= 4p(6D + 8) + (8pD + 23p + m)(24D + 32) \\ &\quad + (2p + a)(Z + J + 12D + 16). \end{aligned}$$

This is the compiler's declared radius bound, not a claim of minimal radius.

[write] These constants, counts and the radius bound restate Report 15's: equations (3) and (18) and the ledger of Section 11; for the universal source they take the values of Section 38 (269,291,358,255 factors). The new content of Report 19 starts with Section 67.

67 Arithmetic access to the ordered factors

Put $r = D + 3 = L - S + 1$, $U = 2D + 6$, $A = 4D + 15$, $V = 2D + 7$, and $C = 4D + 14$. In the following tables, branch e is in moving-branch order and $k = 0, 1$ denotes outbound/inbound. Indexing is zero-based.

E factor	Index	Parameter range
free	$eA + kU$	
behind	$eA + kU + 1 + t - S$	$S \leq t \leq L$
ahead	$eA + kU + r + t - S$	$S + 1 \leq t \leq L$
dispatch	$eA + 4D + 12$	
endpoint	$eA + 4D + 13$	
commit	$eA + 4D + 14$	
direct zero branch j	$pA + j$	$0 \leq j < a$

For P , add $|E|$ to each local index below.

P factor	Local index	Parameter range
phase free	$eC + kV$	
phase near negative	$eC + kV + 1 + t - S$	$S \leq t \leq L$
phase near positive	$eC + kV + 1 + r + t - S$	$S \leq t \leq L$
home phase q	$pC + q$	$0 \leq q < m$

`gate_at(i)` inverts these partitions with integer division and remainder. It constructs one gate with exactly the frozen name, endpoint shapes, class table, key convention, and radius. The three interaction factors occur after both moving-mode travel lists. Direct branches occur after

all moving branches. Home phases occur after all moving-mode phases. The implementation never loops across the travel range merely to locate a factor.

The included small-source differential checks compare complete factor descriptors, rather than only outputs on selected configurations. This tests index boundaries, shape order, contextual domain versus image selection, and radius arithmetic against the original compiler.

68 Complete discovery of possible changes

A candidate set is allowed to contain false positives. Its required property is stronger and simpler than recognition of enabled computation: it must contain every factor with a raw endpoint occurrence.

Lemma 68.1 (Unique close pair within an endpoint). *Every three-particle factor endpoint contains exactly one pair with gap at most D . Its remaining particle lies more than D from each member of that pair. Every two-particle endpoint has its identifying pair.*

Proof. The mode gap lies in $[1, D]$. The singleton's distance from the head anchor is at least $S = 2D + 2$, and therefore its distance from either head particle is at least $S - D = D + 2 > D$. At a translated endpoint reverse shape the marker has moved, but its relative head distance is again S . This proves the assertion about each endpoint; it makes no uniqueness assertion about an entire malformed configuration. $\square$

Consequently, enumerate *every* occupied increasing pair $(h, h + d)$ with $1 \leq d \leq D$. Decode mode and sign from d . For a moving mode enumerate every distinct third occupied site z . Define $H = h - z$ and $\epsilon = 0$ for plus, 1 for minus. Let $w = v$ for outbound and $w = -v$ for inbound. The possible travel parameters follow by inversion:

Family	Recovered parameter	Retain when
behind	$t = wH - \epsilon$	$S \leq t \leq L$
ahead	$t = -wH + \epsilon$	$S + 1 \leq t \leq L$
phase near	$t = H $, side = $\text{sgn}(H)$	$S \leq t \leq L$

The free and phase-free factors are candidates for every detected moving pair. The interaction cases are

Signed mode	Relative position	Candidate
O_e^+	$H = -vS$	endpoint
O_e^-	$H = vS$	dispatch
I_e^+	$H = vS$	commit
I_e^-	$H = -vS$	endpoint

For a home pair the required singleton is $z = h - S$. A plus home pair adds all source-adjacent dispatch and direct factors; a minus home pair adds all target-adjacent commit and direct factors. Either sign adds its home phase. Incidence lists contain $O(b)$ entries in total.

Lemma 68.2 (Candidate completeness). *If a frozen factor g_i has any raw key on finite support X , candidate discovery emits i . Therefore every noncandidate factor is the identity on X .*

Proof. The raw endpoint's close pair is included by the enumeration. In a triple its singleton is also enumerated. The tables above invert every endpoint family, for both signs, directions, and

endpoint sides. The decoded family and parameter recover exactly its index from Section 67. A home endpoint finds the necessary branch through the complete incidence list. Two-particle endpoints directly add the free families. No raw-window, isolation, competitor, context, or machine-admissibility predicate is used to remove candidates. Thus every raw occurrence is covered, regardless of any additional particles in X . Without a raw key the frozen factor has no active substitution. $\square$

Two anchor distinctions are essential. A free reverse endpoint's raw key is the predecessor anchor $h - w$, not its present head anchor h . An endpoint reverse occurrence's invariant key is the new singleton position minus $v\delta$, not the new singleton itself. Candidate discovery only proposes the index; exact shape matching recovers these original keys.

69 Exact sparse application of a factor

For a gate with support bound B , let its two ordered shapes be Q_0, Q_1 . The frozen isolation radius is $\ell = 3B + 1$; its competitor radius is

$$M = \max(\ell + B, T + B) + 1,$$

where $T = 0$ without context and $T = Z + J$ with context. The factor radius is $M + 2B$. The travel constant L from Section 66 and the gate isolation radius ℓ are distinct quantities.

69.1 Raw keys and eligibility

Sort the finite support once. For each shape Q_j and each occupied coordinate f , test the possible raw key $u = f - \min Q_j$. Check membership of all two or three translated shape sites, and count occupied sites in the closed interval $[u - B, u + B]$ by two binary searches. A raw occurrence exists exactly when that count equals $|Q_j|$. Every real raw key occurs in this enumeration because its least shape site is occupied.

Retain *all* raw keys, including those that will fail isolation or context. For sorted distinct raw keys, another key lies at positive distance at most M if and only if an adjacent key in sorted order does. This nearest-neighbor test is therefore exactly the frozen all-raw-key exclusion predicate. In particular, an ineligible raw occurrence can still block another one.

Isolation is the exact inclusive count in $[u - \ell, u + \ell]$. A contextual gate reads the two closed intervals

$$[u - Z - J, u - Z], \quad [u + Z, u + Z + J].$$

An empty read interval represents class $J + 1$; exactly one represents its offset $k \in \{0, \dots, J\}$; more than one rejects that side's guard. Membership in the pinned branch domain or image table then decides the context. Interval counting is equivalent to the original scan over all $J + 1$ sites and does not enumerate that interval.

69.2 Simultaneous substitution

Compute the complete active-key list from the same pre-factor support X . Only afterward replace every translated Q_j by Q_{1-j} . The frozen isolated-swap lemma gives disjoint simultaneous substitutions. In particular, no earlier rewrite within this factor is allowed to affect a later key's eligibility.

Lemma 69.1 (Exact sparse factor). *For every frozen factor g and finite support X , sparse application returns exactly the output of the frozen factor method.*

Proof. Raw shape membership and inclusive interval cardinalities coincide. The adjacent-raw-key test equals the all-keys test. The contextual interval rule returns exactly the original class or rejects exactly the same malformed context. Both algorithms therefore construct the same active substitutions from the same X , and execute the same disjoint replacements. $\square$

Optional verification uses explicit exceptions to check mass, the set of raw keys, and the factor's involution. These checks remain active under Python optimization. Verification is additional executable evidence; it is not an assumption of the main theorem.

70 Ordered skipping and exact inverse

Maintain the current support X and a strict index cursor. A forward call starts below index zero; an inverse call starts at F . Generate candidates on the current support, retain only indices beyond the cursor in the chosen direction, and sort them in that direction.

```

cursor = F if inverse else -1
while True:
    pending = candidates(X) strictly beyond cursor
    sort pending descending if inverse, ascending otherwise
    for i in pending:
        Y = exact_simultaneous_factor(i, X)
        cursor = i
        if Y != X:
            X = Y
            break
            # rediscover on changed support
    else:
        return X

```

Proof of Theorem 65.1. Induct through the original ordered product. By Lemma 68.2, every factor skipped before the next candidate has no raw key and is the identity on the current support. The chosen candidate is executed exactly by Lemma 69.1. If it is the identity, the support and current candidate set remain valid. If it changes support, the cursor advances past it and candidates are regenerated before any further factor is considered. This finds newly created future candidates and reevaluates any changed isolation, competitor, or contextual conditions. Newly created candidates at already-passed indices are correctly ignored because the frozen ordered product does not revisit them. When no candidate remains, the entire remaining suffix acts as the identity. Every tested index advances strictly, so at most F factors are tested and the process terminates. Descending order gives precisely the reversed product. Each factor is its own inverse, so this product is $\mathcal{F}^{-1}$. $\square$

Corollary 70.1. *On all finite supports the lazy evaluator preserves mass, is translation-equivariant, fixes the empty support, and satisfies both inverse roundtrip identities.*

Proof. These properties follow from equality with the frozen ordered composition. They do not require encoded-state admissibility. $\square$

71 A malformed input with a real cascade

Take controls q, h , one true-guard right-counter increment from q to h , and $J = 0$. Then $D = 8$, $S = 18$, $B_3 = 37$, and the outbound signed gaps are 5, 6. On

$$X = \{-118, -112, 0, 18, 23\},$$

the complete forward step changes four factors at global indices 0, 1, 47, 60:

Index	Family	Support after this factor
0	free outbound	$\{-119, -114, 0, 18, 23\}$
1	behind at 18	$\{-119, -114, 0, 19, 25\}$
47	phase free	$\{-119, -113, 0, 19, 25\}$
60	phase near at 19	$\{-119, -113, 0, 19, 24\}$

The first change removes a particle at the inclusive boundary of a later triple’s isolation window, allowing that later factor to act. A one-time candidate or eligibility snapshot would be unsafe.

The same input shows that an entire E or P block need not be an involution on malformed configurations:

$$\begin{aligned} E(E(X)) &= \{-118, -112, 0, 19, 25\} \neq X, \\ P(P(X)) &= \{-118, -112, 0, 18, 24\} \neq X. \end{aligned}$$

This agrees with the frozen proof: block matching is a statement on admissible configurations, while global reversibility follows from the individual involutions and reversed factor order. The evaluator must not substitute a matching shortcut or apply a whole block as its own inverse. The package includes complete event traces for this fixture in both the main and independent audit receipts.

72 Complexity and resource limits

Let G be total parsed guard-program size, $c = (J + 2)^2$, and Δ the largest relevant source/target incidence list for a home signed mode. Let K be the number of changing factors in one step, A_{test} the number tested, and C_X an upper bound on candidate-set size across its snapshots. Then

$$K \leq A_{\text{test}} \leq F, \quad \text{candidate generations} = K + 1, \quad C_X \leq F, \quad C_X = O(n^3 + n^2\Delta).$$

Mass remains n throughout.

72.1 Source metadata and class cutoff

Metadata occupies $O(m + b + G + bc)$ words, plus label strings and numeric payloads. This includes the immutable source snapshot. With unit-cost integer and hash operations, class-table work is $O(cG + bc)$. The actual growing integer-mask implementation performs c growing-mask operations per branch. A deliberately loose bit-work bound for those operations is $O(bc^2)$, in addition to guard evaluation and integer comparisons. Thus a large binary-encoded J remains expensive. No polynomial-in-log J validation bound is claimed.

The source-sized characterization refers to source records and their explicit class tables, rather than F factors. It does not claim that an arbitrary source can be validated without reading its explicit input, or that class tables are small for every cutoff.

72.2 Expected hash model

With customary constant-expected-time hashing and unit-cost arithmetic, one step excluding metadata and optional event storage costs

$$O\left((K+1)[n \log(n+1) + n^3 + n^2 \Delta + C_X \log(C_X + 1)] + A_{\text{test}} n \log(n+1)\right).$$

Candidate enumeration and sorting give the first term; sparse gate application gives the second. Actual gate-name string-copy costs are additional. Work memory is $O(n + C_X)$ beyond fixed metadata. With `trace=True`, active-key event storage can add $O(Kn)$ coordinate words; with `trace=False` these events are not retained.

72.3 Adversarial hashing and exact integers

Python integer operations and hash tables do not provide constant-time worst-case primitives. Charging candidate insertions and dictionary lookups by linear container sizes gives the following conservative container-operation bound:

$$O\left((K+1)[(n^3 + n^2 \Delta)(C_X + n + m + b + 1) + C_X \log(C_X + 1)] + A_{\text{test}}[n^2 + n \log(n+1)]\right).$$

It is intentionally loose. Source-validation dictionaries have analogous collision costs. Adversarial integer hashes affect speed, not the exactness argument.

Let H_0 bound initial absolute coordinates. Every endpoint lies within $[-B, B]$ with $B \leq B_3$, and its paired endpoint has equal cardinality. Pairing replaced particles in sorted order therefore bounds each displacement by $2B_3$. After K changing factors,

$$|x| \leq H_0 + 2KB_3$$

for every occupied coordinate. Context arithmetic extends at most $Z + J$ farther. A uniform coordinate bit bound is

$$W = O(1 + \log(1 + H_0 + 2FB_3 + Z + J)).$$

Charge coordinate addition, comparison, and hashing by $O(W)$ in a simple bit model; factor indices require $O(\log(F+1))$ bits. Other source integers and label payloads carry their actual costs. The coordinate span itself is never scanned.

No uniform short cascade. The worst case retains K and A_{test} , potentially as large as F . Malformed support can create future eligibility through geometry, contexts, isolation, or raw competitors. Although $F = O((m+b)^2)$, avoiding allocation of F factors does not establish an input-independent bound polynomial in $\log F$. Few tested factors in a benchmark are observations about those inputs, not a theorem about all supports.

73 Executable evidence and universal observations

73.1 Differential regression and independent audit

Every count below is per execution mode. Normal Python and `python -O` rerun the same cases; they are not counted as distinct mathematical inputs.

Main regression	Count
Small accepted sources	12
Complete factor descriptor comparisons	3,580
Whole-step lazy versus eager comparisons	17,517
Sparse individual-factor comparisons	71,695
Complete candidate-completeness sweeps	241
Invalid-input and immutable-metadata checks	35
Independent audit	Count
Small source fixtures	5
Complete factor descriptor comparisons	1,471
Raw endpoint candidate checks	14,710
Sparse individual-factor comparisons	52,956
Contextual guard checks	268
Shape/random whole-step comparisons	830
Exhaustive support/order comparisons	8,192
Validation acceptance and exact-error fuzz cases	1,000
Typed invalid cases	15

The 8,192 exhaustive comparisons are the 2^{11} subsets of each of two eleven-site universes, each tested in both orders: 4,096 distinct supports. With the 830 shape/random comparisons, the independent audit performs 9,022 whole-step comparisons per mode. Whole-step cases include inverse roundtrips and mass checks. Both suites share the same pinned reference by design; their separately written comparison paths and the mathematical audit are not claims of formally independent foundations.

The focused addendum adds 123 source-outcome comparisons, including 32 accepted sources at $J = 17,64$ and 54 malformed-schema cases; 864 contextual patterns; 2,672 gate-bound checks; and 23,488 endpoint-pair displacement checks. Its uniform validation argument assumes enough resources for both validators to finish; it does not claim identical resource-exhaustion behavior. The 121 packaging checks cover exact container and element types, Boolean flags, direct-constructor rejection, nested immutability, snapshot independence, empty support, and rejection of a tampered compiler before execution. Detailed counts and precise tested hashes are in the accompanying receipts, rather than inferred from this summary.

73.2 Literal universal benchmark

The supplied universal source has the following exact ledger:

Quantity	Value
Controls m	122,622
Moving branches p	66,066
Zero branches a	75,495
Total branches b	141,561
Class cutoff J	0
D	509,508
S	1,019,018
B_3	2,038,037
Z	20,380,380
$ E $	134,645,688,597
$ P $	134,645,669,658
F	269,291,358,255
Declared radius bound R	3,292,955,588,459,274,804

The source counts are also independently recomputed from the pinned table. The benchmark starts at the published support

$$\{-20,380,381, 0, 1,019,018, 1,019,019, 20,380,380\},$$

which is $\text{enc}(\text{START}, 1, 0, +)$. It evaluates 128 actual CA microsteps. Every step uses the global finite-support evaluator, checks its reverse, and preserves five particles. All outputs are additionally checked against explicit startup geometry. For zero-based output index $i < 3$, the expected home controls are `h0000B0T1`, `p00001R0T1`, and `p00001R0T2`. For $3 \leq i \leq 127$, the fixed markers are $\{-Z - 1, 0, Z\}$ and the plus outbound head is anchored at $-S - (i - 3)$ with gap $2m + 1$. The recorded final support is $\{-20,380,381, -1,019,142, -773,897, 0, 20,380,380\}$. This independent expected-output check supplements the inverse checks; it does not extend the observation window. It then evaluates 128 deterministic arbitrary endpoint/noise cases spread across the full factor-index range, alternating direction and checking inverse and mass. One translated input uses a 512-bit-scale coordinate. No source-boundary jump replaces any of these calls.

The producer's final normal receipt records 3.811 seconds for metadata compilation, 0.0412 seconds for the startup steps with their checks, and 0.0373 seconds for the malformed cases with their checks. Process peak resident size was 376,048 KiB (about 367.2 MiB), including parsed source, metadata, and traces. The largest tested/changed factor counts were 4/2 on startup and 8/2 on malformed cases. The optimized run also passed; its compilation observation was 5.461 seconds under concurrent load. These are environment-specific observations. Fresh portable replay timings are separately recorded and need not match.

What was not completed. The 128 startup microsteps do not complete startup, one Turing transition, or a halting computation. No eager universal factor-array comparison was attempted. The sample demonstrates literal global evaluation without mandatory factor materialization; all-finite-input equality is supplied by the proof and small-source differential comparisons.

74 Public API and portable reproduction

74.1 Exact types and immutable results

The supported API is Python 3.12 or later with the standard library only. `positions` must have exact type `set` or `frozenset`, and every element must have exact type `int`. Booleans, floats, container subclasses, lists, tuples, and arbitrary iterables are rejected. Thus no coercion

silently changes the declared finite binary input. Empty support is accepted. Integer magnitude is unrestricted by the API.

Flags `inverse`, `verify`, and `trace` require exact Booleans. Factor indices require exact nonnegative integers below F . Controls and signs require exact strings; encoder counters require exact nonnegative integers. Direct construction of private frozen records is rejected.

`step` returns a `frozenset`. With tracing it returns the support and an immutable mapping containing tested-factor count, changed-factor count, generation count, and a tuple of event records. The source and branch snapshots, incidence maps, and control map are immutable. `ledger()` returns a detached ordinary dictionary; modifying it does not mutate the rule. `candidate_indices(X)` returns sorted current candidates, and `gate_at(i)` materializes one original factor.

```
import json
from lazy_reversible import compile_lazy_source

with open("source.json", encoding="utf-8") as handle:
    rule = compile_lazy_source(json.load(handle))
x = rule.encode(rule.start, 1, 0)
y, stats = rule.step(x, verify=True, trace=True)
if rule.step(y, inverse=True, verify=True) != x:
    raise RuntimeError("roundtrip failed")
```

74.2 Package contents and replay commands

The ZIP contains this PDF and editable LaTeX, build script, exact evaluator and reference compiler, exact compressed source, the proof and audit narrative, all test scripts, the explicit cascade fixture, original normal/optimized receipts and traces, portable replay receipts, and file manifests. Standalone copies of Reports 15 and 17 provide background; neither external sibling directories nor network access are required.

From the extracted top-level directory run:

```
python3 verify_release.py
python3 release_checks.py
python3 replay.py --out replay-output
```

The verifier checks the manifest and semantic pins. Replay verifies the package first, decompresses the source into a fresh working directory, checks its uncompressed SHA-256, and runs the main suite, independent audit, addendum, packaging checks, and universal benchmark in both ordinary and optimized Python. It compares deterministic counts, checks receipt linkage, and writes new logs and receipts under the chosen output directory. Original evidence is left untouched. All test failures use explicit exceptions or nonzero exit checks, so `-0` does not remove them.

The evaluator requires only the Python standard library. The supplied full replay was tested on Python 3.12.14 on Linux and targets Python 3.12 or later on Linux because the preserved benchmark uses the Unix resource module and Linux KiB RSS units. PDF rebuilding separately requires a working TeX Live installation with the packages named in `report19.tex`; run `sh build.sh`. Rebuilding the PDF is not needed to reproduce any mathematical or computational check. Replayed timings and process resource figures are not byte-deterministic; pinned sources, core code, expected counts, cascade outputs, and universal support traces are deterministic.

74.3 Adaptation and integrity boundary

The evaluator and frozen compiler are copied byte-for-byte. The source is compressed losslessly with a deterministic gzip header. The only benchmark adaptation replaces its environment-specific absolute source path by an adjacent `source.json`; replay places the verified decompressed bytes there. Its sampling seed, algorithms, step counts, and source hash are unchanged. The independent audit and its addendum are byte-preserved. Producer evidence remains distinguishable from fresh portable evidence.

The manifest records every delivered file’s digest except itself and its detached digest. Verification rejects unexpected files and directories outside the exact generated trees `.build/` and `replay-output/`, and rejects symlinks. A custom replay output must lie outside the release; the default output directory must not already exist. This detects accidental or unilateral file corruption; a digest supplied together with its manifest is not a digital signature or an external trust anchor. The hard-coded compiler and source pins provide a reproducible identity boundary. The evaluator checks its adjacent compiler before executing it; changing that dependency causes an immediate failure.

75 Conclusions

Arithmetic factor access and raw-shape candidate completeness remove the requirement to allocate the frozen ordered product. Exact simultaneous gate application and strict-cursor rediscovery preserve its behavior on every finite support, including malformed cascades. Reverse evaluation is the true reversed factor sequence. The gain is a practical finite-support route to this already-defined rule with explicit source- and cascade-sensitive costs. The universal sample establishes literal microstep execution only; startup completion, full Turing transitions, uniform short cascades, and new universality or priority claims are outside this report.

76 Byte identities

[write] This section is Report 19’s Appendix A. The following SHA-256 identities are complete, not abbreviations.

Frozen compiler

f95a6028c9d3cc6f8b0b497cfc0cad68104da6d70e55205df04ad234e5d0c76f

Lazy evaluator

42e8aa65c05fcf373a03a02be51ebb89a4fdbcb1f1070ad776ffe1e8e99049c61

Uncompressed universal source

fa61d06178d718511232a80a05e02192d940fde3c273f0c202f0623253b2ade3

The source contains exactly 32,034,272 bytes before gzip compression. The package’s provenance record also pins the original and adapted benchmark scripts, reference PDFs, and compressed source. Full release identities are in `release-manifest.json` and its detached digest.

Part V

Event-budgeted quartic certificates for the ordered rule (Report 20)

Event budgeted quartic certificates for exact finite CA iteration

Canonical natural witnesses and fully charged finite schemas

Report 20 Mathematical construction and portable reproducibility package

3 October 2026

Abstract of Report 20, as delivered. Fix a source for the frozen reversible binary cellular automaton, an exact particle count n , a finite horizon T , and a total budget K for changing local factors. We construct a finite sum of squares of quadratic residuals whose natural auxiliary fiber is a singleton precisely when the literal ordered CA product completes T steps within that budget and meets a prescribed endpoint, if any. The construction uses $T + K$ deterministic rounds: changing-factor rounds, one separate absence-and-completion round per CA step, and forced idle padding. It certifies simultaneous local-factor semantics and the identity of every omitted suffix. A support-indexed candidate list replaces an array indexed by every factor; source records, finite guard tables, comparisons, sorting, quotient/remainder, dummy wires, and full polynomial expansion remain charged. We prove complete auxiliary uniqueness by a deterministic primitive circuit. A general fixed-slot scheduler and an implemented mass-two coefficient exporter provide reproducible tests, including a complete 612-witness, 616-residual sample and its fully expanded quartic. The arbitrary-mass coefficient exporter is not implemented. This is a fixed-parameter finite-schema theorem, with no real-witness, unknown-horizon, finite-fold MRDP, mass-two universality, novelty, or optimality claim.

[write] Delivered as `Event_Budgeted_Quartic_Certificates_Package.zip` (batch 82, archive 08; arrival db37d18c8), dated 3 October 2026; files placed by 7d2b1b245 with the prefix `20-event-budget-` (Appendix A). Report 20 certifies the frozen *ordered* rule of Report 15’s compiler (Part I), whose exact finite-support evaluation is Report 19’s Theorem 65.1 (Part IV); its universal instance is Report 17’s source (Part III). Its arithmetic gadgets are those of “Report 8”, which is source 13 (Part IV) of `signal-machine-collision-certificates` (SMC); the bundled excerpts `report8-core.tex` and `report8-semilinearity.tex` are SMC’s own text and are not shipped, and the circuit-to-quartic layer is credited to `canonical-diophantine-certificates` (CDC) in the note after Lemma 79.1. Report 28 (Part VI) uses the same compiler pattern for a different rule and a different theorem (note in Section 114). The Hilbert’s-tenth programme’s scoped intake `Computability/HilbertTenthProblem/Papers/research-wip/native-stream-queue/review_event_budget_clean_clock_intake.md` (commit 8aabb1453) read Report 20’s complete proof without executing its Python, identified no new defect within its stated theorem scope, and records that the witness arity grows with n , T , K and the source, so the construction gives no fixed-arity representation and no ordinary-input loader for Report 16’s source (Section 1.6). The delivered table of contents is omitted, and Report 20’s H_* is printed Ht_* (Section 85).

77 Result and implementation boundary

A finite binary configuration is represented by a strictly increasing tuple $X = (x_1, \dots, x_n)$ of integer occupied sites, with zeros elsewhere. Coordinates are unbounded in both directions. Fix

a source accepted by the frozen compiler of the earlier construction. Its ordered factors are $g_0, \dots, g_{F-1}$ and one literal CA step is

$$\mathcal{F}(X) = g_{F-1}(\cdots g_1(g_0(X)) \cdots).$$

Report 19 proves that the pinned lazy evaluator computes this exact product on every finite support, including malformed configurations [17]. This report changes the certificate representation, not the rule or its factor order.

A *changing factor* is an application of g_i whose simultaneous output differs from the entire pre-factor support. Several separated active occurrences of one factor count as one changing factor. Let $k_T(X)$ be the total number of such applications during the literal evaluation of T CA steps. This number is determined by the input and the ordered product; it is not a guessed event count.

Theorem 77.1 (Finite event-budgeted canonical natural fiber). *Fix an accepted finite source, exact mass n , horizon $T \in \mathbb{N}$, and budget $K \in \mathbb{N}$. Initial coordinates are supplied as $2n$ natural inputs*

$$x_i = x_i^+ - x_i^-, \quad x_i^+ x_i^- = 0, \quad x_1 < \cdots < x_n.$$

There is an effectively generated integer polynomial

$$P(X^+, X^-, w) = \sum_{j=1}^E \rho_j(X^+, X^-, w)^2, \quad \deg \rho_j \leq 2,$$

with the following properties for every such fixed admissible input.

1. *If $k_T(X) \leq K$, its natural auxiliary fiber contains exactly one point. With a prescribed sorted endpoint Y , that point exists exactly when also $\mathcal{F}^T(X) = Y$.*
2. *Otherwise the natural auxiliary fiber is empty.*
3. *The witness determines every circuit wire, all tested snapshots and factor effects, each actual changing factor, each step-completion absence check, and all idle padding.*
4. *No auxiliary array or lookup table is indexed by all F factors. The actual source and guard-table costs are retained explicitly.*

The construction is finite for each fixed (n, T, K) and source. Its quantified dimension depends on these parameters.

The external input variables are not counted as existential witnesses. Canonical signs and initial strict order can instead be checked internally with the additional ledger below. For $T = 0$, the externally promised canonical sorted domain admits a witness-free endpoint polynomial. Internal rejection of arbitrary unsorted tuples requires comparison auxiliaries, so that stronger interface is not witness-free. For $n = 0, 1$, every factor is the identity and a witness-free identity specialization is also available. The uniform scheduler works at these masses with zero candidate slots.

What is implemented. The package supplies a general fixed-slot scheduler reference and tests against the pinned factor evaluator. It separately supplies a complete coefficient emitter for all two-particle inputs of one small fixed source. The arbitrary-mass theorem is given by the explicit finite circuit construction in this report; there is no shipped arbitrary-mass coefficient exporter. The conservative macro bound is a theorem-level upper bound, not a measured general compiler count. The small sample's coefficient counts are mechanically emitted and checked.

78 Source data and inherited local semantics

78.1 Constants and costs

Write m for the number of controls, p for moving branches, a for direct branches, and $b = p + a$. The source has a nonnegative guard cutoff J . The unchanged constants are

$$\begin{aligned} D &= 2m + 4p, & S_0 &= 2D + 2, & L_0 &= 3D + 4, \\ B_3 &= 4D + 5, & Z &= 10B_3 + 10 + 2J, & F &= 8pD + 29p + m + a. \end{aligned}$$

Let Δ be the maximum size of a signed home-mode incidence list. Its outgoing entries are dispatch/direct factors and its incoming entries are commit/direct factors. There are exactly $2b$ actual incidence records in total. Set

$$c = (J + 2)^2, \quad S_{\text{tab}} = m + b + 2bc + 1.$$

The $2bc$ term counts the actual domain and image guard-table entries. Source names, guard syntax, and their byte lengths are separate setup data. No F -entry table, packed F -bit truth table, or cost-free giant coefficient is admitted.

78.2 What one factor means

A local-factor descriptor contains two equal-size endpoint shapes of size two or three; a local isolation parameter B ; the isolation radius $3B + 1$; a competing-key radius

$$\max(4B + 1, T_{\text{context}} + B) + 1;$$

and, when contextual, the branch and domain/image guard identifier. Here T_{context} is a descriptor radius, unrelated to the CA horizon T . Endpoint shapes are not translates of each other. A raw key identifies a translated endpoint that occurs with the exact local particle count. An active key additionally passes isolation, the all-raw-key competitor test, and its context guard.

The inherited isolated-swap theorem says that active endpoint unions are disjoint and the factor simultaneously swaps all active endpoints. A factor is never a sequential scan that rewrites one occurrence and then decides another on the modified configuration. The complete descriptor definitions, arithmetic indexing proof, and all-finite-support completeness lemma are included as the bundled Report 19 proof [17]. They are mathematical dependencies, not assumptions about well-formed encoded inputs.

78.3 Arithmetic factor indices

For clarity, here are the fixed formulas used by the schema. Put

$$r = D + 3, \quad U_E = 2D + 6, \quad A_E = 4D + 15, \quad V_P = 2D + 7, \quad C_P = 4D + 14,$$

and $E_0 = pA_E + a$. A moving branch index is $e \in \{0, \dots, p - 1\}$; $k = 0, 1$ denotes out-bound/inbound. Indices are zero based.

Factor	Index	Parameter range
E free	$eA_E + kU_E$	
E behind	$eA_E + kU_E + 1 + t - S_0$	$S_0 \leq t \leq L_0$
E ahead	$eA_E + kU_E + r + t - S_0$	$S_0 + 1 \leq t \leq L_0$
Dispatch	$eA_E + 4D + 12$	
Endpoint	$eA_E + 4D + 13$	
Commit	$eA_E + 4D + 14$	
Direct branch j	$pA_E + j$	$0 \leq j < a$

The P factors follow the E block. Home phase for control q has index $E_0 + pC_P + q$. For moving branch e , the free phase has index $E_0 + eC_P + kV_P$. The two near-phase families have indices

$$\begin{aligned} \text{negative side: } & E_0 + eC_P + kV_P + 1 + t - S_0, \\ \text{positive side: } & E_0 + eC_P + kV_P + D + 4 + t - S_0, \end{aligned}$$

where $S_0 \leq t \leq L_0$ and the side is the sign of $H = h - z$. These formulas invert into a descriptor using comparisons, fixed-divisor quotient/remainder, and source-record selection. They do not require enumerating F factors.

79 Canonical quadratic primitive compiler

Every internal signed wire v is represented by two natural variables v^+, v^- . An arithmetic assignment $v = f$ emits

$$v^+ - v^- - f = 0, \quad v^+ v^- = 0. \quad (79.1)$$

For any fixed preceding values there is exactly one natural solution: the positive and negative parts of f . Allowed assignments are binary addition, subtraction, multiplication, fixed-integer scaling, and selection

$$f = u + e(w - u),$$

where e is a previously determined Boolean wire. Long sums become binary folds. Products of already quadratic expressions are materialized first; all defining residuals have degree at most two.

Lemma 79.1 (Unique comparison and constant division). *For fixed signed integers x, y , natural β, d satisfying*

$$\beta(\beta - 1) = 0, \quad y - x - (2\beta - 1)d - \beta + 1 = 0 \quad (79.2)$$

are unique. They are $(1, y - x)$ if $y \geq x$ and $(0, x - y - 1)$ otherwise. For a fixed positive integer d_0 , natural q^+, q^-, r, h satisfying

$$q^+ q^- = 0, \quad v - d_0(q^+ - q^-) - r = 0, \quad r + h - (d_0 - 1) = 0 \quad (79.3)$$

are unique and give the Euclidean quotient and remainder of v by d_0 .

Proof. The first residual in (79.2) restricts β to 0 or 1. In each case the second fixes d , and naturality admits precisely the indicated case. In (79.3), naturality gives $0 \leq r < d_0$. Euclidean quotient and remainder are unique even when v is negative; sign separation and the last equation then uniquely fix all four variables. $\square$

[write] **Credit and a second presentation.** The circuit-to-quartic layer of this section and of Section 83 is already printed in the collection. Signed natural pairs (79.1) are Lemma `cdc:pt:lem:signed`, and the topological induction with the sum of squares is the canonical circuit lemma `cdc:pt:lem:circuit`, of `canonical-diophantine-certificates` (CDC); the comparison gadget (79.2) is CDC's Lemma `cdc:mem:lem:compare`, printed also in SMC Part IV as `smc:s1:eq:cmp`; equality as a difference of two comparison bits is SMC's `smc:s1:eq:eq`; the constant division (79.3) is the congruence gadget of SMC's Section `smc:s1:sec:presburger-quartic`. Report 20 cites Report 8 (SMC source 13) for these gadgets and proves them again; it does not cite CDC. The section is kept as Report 20's own presentation; neither Report 20 nor this report claims novelty for it.

Equality is built from the flags of $y \geq x$ and $y \geq x + 1$, with their difference materialized as a Boolean output. The $x + 1$ assignment is charged. Ties have a prescribed result. Absolute value

is a comparison plus signed selection; order, minimum, floor, and remainder are never implicit primitives.

Boolean AND, OR, NOT, and affine copy use one natural output z and respectively the residual

$$z - uv, \quad z - u - v + uv, \quad z - 1 + u,$$

or the appropriate affine equation. Their inputs are already forced Boolean wires, so Booleanity and uniqueness follow inductively. All large conjunctions and disjunctions are binary trees or folds.

79.1 Exact primitive ledger

Let A, I, D_v, G be the numbers of arithmetic assignments, comparisons, constant divisions, and Boolean gates. Let Q count acceptance and external-admissibility residuals. Then

$$V = 2A + 2I + 4D_v + G, \tag{79.4}$$

$$E = 2A + 2I + 3D_v + G + Q. \tag{79.5}$$

These are exact compiler ledgers, before any optional elimination of wires or residuals. The following coefficient-list bounds retain zero or cancelling monomial slots if necessary:

$$L_1 \leq 9A + 10I + 9D_v + 4G + 4Q,$$

$$L_2 \leq 65A + 68I + 35D_v + 16G + 16Q.$$

Here L_1 is the number of residual monomial slots and L_2 is the number of ordered term occurrences in all expanded squares. Assignment row lengths are at most 8, 1; comparison row lengths at most 2, 8; division row lengths at most 1, 5, 3; Boolean rows and final checks at most four. Squaring each row gives these stated bounds. Collected expanded monomials may be fewer, but their collection is a separately charged operation.

False masks do not erase wires. Every internal gate below a false mask is still evaluated and constrained. An exposed inactive candidate index is F ; unused shape positions are zero; idle states retain their previous values. Internal dummy wires have unique canonical values, which need not be zero. Merely multiplying a defining equation by an enable bit would leave hidden freedom and would not prove this theorem.

80 Fixed candidate slots without a factor array

For every occupied pair $x_i < x_j$, let $d = x_j - x_i$. Allocate the following ordered slots:

1. one home-phase slot;
2. Δ home-incidence slots in the pinned incidence order;
3. one moving free- E slot and one moving free- P slot;
4. for every $k \notin \{i, j\}$, four slots: behind, ahead, phase-near, and interaction.

Thus there are exactly

$$C = \binom{n}{2} [\Delta + 3 + 4 \max(n - 2, 0)] \tag{80.1}$$

slots. A slot consists of an enable bit and an index; disabled slots carry index F . Duplicates are retained, with fixed ranks. Only the enabled set, not an arbitrary ordering of that set, has to match the candidate set of Report 19.

For a home gap $1 \leq d \leq 2m$, decode

$$q = \left\lfloor \frac{d-1}{2} \right\rfloor,$$

with odd/even gap giving plus/minus. A home phase or incidence slot is enabled only if $x_i - S_0$ is occupied. The incidence records are ordered by the source-order-preserving moving-branch filter followed by the source-order-preserving direct-branch filter. This need not agree with the original interleaved branch order.

For $2m < d \leq D$, quotient/remainder of $d - 2m - 1$ by 4 gives the moving branch, outbound/inbound, and sign. Put $h = x_i$, $z = x_k$, $H = h - z$, and $\epsilon = 0$ for odd d and 1 for even d . With branch side s , set $w = s$ outbound and $w = -s$ inbound. Both free slots are enabled. The triple slots test

$$\begin{aligned} \text{behind: } S_0 &\leq wH - \epsilon \leq L_0, \\ \text{ahead: } S_0 + 1 &\leq -wH + \epsilon \leq L_0, \\ \text{phase near: } S_0 &\leq |H| \leq L_0. \end{aligned}$$

The interaction slot uses four mutually exclusive cases:

Mode	Additional equality	Factor
Outbound plus	$H = -sS_0$	Endpoint
Outbound minus	$H = sS_0$	Dispatch
Inbound plus	$H = sS_0$	Commit
Inbound minus	$H = -sS_0$	Endpoint

All factor indices use the formulas above. A disabled mode is masked to a valid dummy before division, record lookup, or range-dependent arithmetic. If $p = 0$, use a constant dummy moving record with side +1 and false enable, rather than selecting a nonexistent record. A disabled travel parameter is masked to S_0 .

80.1 Paying for source selection

A table of size $m\Delta$ is unnecessary and not hidden in the cost. Each of the $2b$ actual incidence records stores its control, sign, rank, and factor index. Compare the selected control/sign with each actual record and route its masked contribution to its fixed rank accumulator. Each rank contains its unique match or the sentinel. Comparisons, masks, and accumulator additions are all circuit operations. Moving-record attributes use a linear-size multiplexer over the actual p records.

Lemma 80.1 (Candidate completeness). *Every nonidentity factor on the current support appears in at least one enabled slot. A factor outside the enabled set has no raw occurrence.*

Proof. Each two-particle endpoint has an occupied close pair. Every triple endpoint has exactly one close pair; the third particle is one of the enumerated x_k . The gap decoder and the displayed finite cases are exactly those of the pinned all-finite-support candidate lemma. Hence each raw occurrence yields a slot for its factor. Isolation, guards, and well-formed encoding are not used in this enumeration, so malformed supports are included. Repeated discoveries merely yield duplicate slots. $\square$

81 Exact simultaneous evaluation of each candidate

The scheduler evaluates the entire effect of every candidate from the same current snapshot. A disabled slot selects factor zero as a total dummy and masks its final effect. A valid source has $F \geq 1$. The descriptor is decoded arithmetically, with source-linear multiplexers. It contains two endpoint shapes of equal size $s \in \{2, 3\}$ in fixed three-position arrays. Sort by presence first and coordinate second, so the unused zero record is last even if actual shape coordinates are negative. A fixed three-record compare-exchange network suffices.

81.1 Raw keys and isolation

For label $\ell \in \{0, 1\}$ and each occupied coordinate x_i , form

$$u_{\ell,i} = x_i - \text{shape}_\ell[0].$$

There are exactly $2n$ key slots. The raw flag requires every necessary translated shape site to be occupied, using equality against all n coordinates. It also requires exactly s particles in the inclusive interval $[u - B, u + B]$, using two comparisons per particle and a binary count sum.

Within one label, distinct x_i yield distinct keys. Across labels, the nontranslation property prevents two raw labels at one key. Thus raw keys have no duplicate valid representations; no existential choice of occurrence is present.

Isolation uses the same inclusive counting circuit on $[u - (3B + 1), u + (3B + 1)]$. Every key slot is compared against all $2n$ raw slots and is rejected when a *different raw key* lies within the competing-key radius. Competitors include raw keys that fail their own isolation or context. Discarding such failed competitors would change the factor and is unsound.

81.2 Finite context guards

For each contextual side, count occupied sites in

$$[u - Z - J, u - Z] \quad \text{or} \quad [u + Z, u + Z + J].$$

A count greater than one rejects the guard. Zero gives the overflow class $J + 1$. One gives the exact offset of its unique particle, obtained by a stable first-hit scan with Boolean selections. On a rejected multiple-particle context, select a valid dummy class before table lookup.

The lookup compares the two class fields and the branch/domain-image identifier *fieldwise* with all actual $2bc$ table entries, and forms the uniquely selected output by a finite masked sum. A flattened address such as $c_0(J + 2) + c_1$ is deliberately not formed: it would introduce an unnecessary J^2 -scale wire. If $b = 0$, the unused lookup has a fixed true dummy output and accesses no nonexistent row. No guard routine scans a physical interval of $J + 1$ coordinates.

81.3 The simultaneous output and canonical order

All active flags are decided from the unchanged pre-factor snapshot. Pair the sorted old endpoint sites with the sorted new endpoint sites by rank. For each old particle x_j , construct

$$y_j = x_j + \sum_{a,r} [\text{active}_a \wedge x_j = u_a + \text{old}_{a,r}] (\text{new}_{a,r} - \text{old}_{a,r}). \quad (81.1)$$

The rank r ranges over live shape positions; the padded third position has zero displacement. The sum has at most $6n$ terms for each j , expanded into binary additions and lifted products.

By the isolated-swap lemma, active endpoint unions are disjoint, so at most one displacement applies to any old particle. Formula (81.1) is exactly the simultaneous factor output, without particle loss, ambiguous labels, or a serial-rewrite artifact.

Sort the n output coordinates with a fixed adjacent compare-exchange network using $\binom{n}{2}$ comparators, keeping the earlier record on ties. Factor semantics ensure distinct output coordinates. The sorted output is therefore canonical. Compare it coordinatewise with the input to obtain the changing bit. This comparison avoids needing to assume that a nonempty active list must change the set, even though that implication is available from the swap construction.

82 Least future change and the completion rounds

For $T > 0$, set $R = T + K$. The state at round r is a sorted support X_r , a step counter t_r , and a cursor c_r denoting the least unprocessed factor index. Initially $t_0 = c_0 = 0$. In every round, build all C slots and all their factor outputs from X_r .

Scan the slots in their fixed order. Start with best index F and retain a candidate precisely when

$$\text{enabled, } i \geq c_r, \quad g_i(X_r) \neq X_r, \quad t_r < T,$$

and its index is strictly smaller than the current best. The stored output is selected coordinatewise. Strict replacement forces the earliest slot among duplicated minimum indices. Comparisons and selections explicitly implement this minimum; there is no minimum oracle.

There are three mutually exclusive updates.

1. **A winner exists.** Apply its already computed output; set $c_{r+1} = i + 1$ and keep $t_{r+1} = t_r$. Recompute the candidate list from the new snapshot next round.
2. **The step is unfinished and no winner exists.** Keep the support, set $t_{r+1} = t_r + 1$, and reset $c_{r+1} = 0$.
3. **The horizon is complete.** Keep the support and $t = T$, with cursor zero. Every eligibility bit is false; all internal computations still have uniquely forced values.

The final acceptance residual requires $t_R = T$, along with any endpoint equations.

Lemma 82.1 (Absence certifies the remaining ordered suffix). *In the second case, every remaining factor is the identity on the current snapshot, so the update completes exactly one literal CA step.*

Proof. Every future candidate was evaluated on that snapshot and failed to change it. By Lemma 80.1, every noncandidate factor has no raw occurrence. Each factor in the remaining ordered suffix is therefore the identity. Applying these identities in their literal order leaves the same snapshot for the next factor, so the full suffix is the identity. Testing only a proposed winning event, without these absence decisions, would not establish this fact. $\square$

Lemma 82.2 (Exact total budget). *An execution with k changing factors over T complete steps uses exactly $k + T$ nonidle rounds, followed by uniquely forced idle padding. It reaches $t_R = T$ by $R = T + K$ if and only if $k \leq K$.*

Proof. Induct on the literal ordered factor sequence. In a changing round, the minimum future changing candidate is the next nonidentity factor: preceding factors are identities, so they preserve the snapshot on which that choice was tested. After a change, recomputation discovers any newly created later candidates while never revisiting passed indices. By Lemma 82.1, a no-winner round completes exactly the remaining suffix of one step. There is one such round per

step, including the step containing the last allowed change. Thus there are exactly k changing rounds and T completion rounds. All later rounds are idle, proving the equivalence. $\square$

The $+T$ is indispensable: K event slots alone cannot pay for the final absence check of every step. Conversely, duplicating an event or choosing a later eligible factor cannot manufacture a new witness, because the stable minimum circuit fixes the schedule.

83 Proof of the complete natural fiber statement

For a fixed admissible input, evaluate the straight-line circuit in its declared order. At each arithmetic assignment, equation (79.1) has one natural solution. At each comparison or constant division, Lemma 79.1 gives one natural extension. Boolean gates have uniquely determined outputs, inductively Boolean. Finite multiplexer sums, sorting networks, minimum scans, and scheduler updates are compositions of these primitives. All masked internal gates remain constrained and totalized.

It follows by induction over *every primitive gate*, not merely over the displayed state snapshots, that the pre-acceptance system has exactly one complete natural assignment. This includes all inactive candidate descriptors, guard dummies, quotient/remainder slacks, losing outputs, and idle-round computations. The acceptance residuals either retain that assignment or remove it. Lemma 82.2 identifies completion with $k_T(X) \leq K$, and the target equations identify its output with the prescribed endpoint.

Finally, over natural numbers (indeed over integers), a sum of integer squares is zero exactly when every residual is zero. The sum-of-squares polynomial has degree at most four and exactly the same natural fiber as the residual system. This proves Theorem 77.1.

[write] The induction above is CDC's canonical circuit lemma `cdc:pt:lem:circuit` (note after Lemma 79.1). The closest theorem in the collection is SMC Part IV's Theorem `smc:sl:thm:core`, a sparse finite-horizon quartic with exactly one natural zero for number-conserving partitioned automata. Theorem 77.1 is a statement of that kind for a different automaton, the frozen ordered rule; its new ingredients are the event budget K , the T separate absence-and-completion rounds and the occupied-pair candidate slots, which make the circuit scale with $(T + K)C$ rather than with T times the factor count F .

The argument is specific to natural auxiliaries. The comparison slack, sign separation, and nonnegativity conditions are not a claim of unique or equivalent real solutions. Finite one-coordinate witness mutation tests support the implementation, but they do not replace this deterministic-induction proof of complete uniqueness.

84 Fully charged size and expansion bounds

For $T > 0$, define

$$U = 10000(T + K)(C + 1)(n + 1)^2 S_{\text{tab}}. \quad (84.1)$$

The binary-fanin construction can be chosen so that

$$A + I + D_v + G \leq U.$$

This is intentionally generous, with all source multiplexer and table costs retained. The per-round macro allocation below makes its scale explicit.

Work	Primitive upper allocation
Candidates, incidence selection, and all index/descriptor decoders	$1000(C+1)(n+1)S_{\text{tab}}$
All candidate effects, raw and competitor tests, contexts, table selection, displacement, sorting	$1000(C+1)(n+1)^2S_{\text{tab}}$
Stable minimum, completion and idle flags, state updates	$1000(C+1)(n+1)S_{\text{tab}}$

These sum to at most $3000(C+1)(n+1)^2S_{\text{tab}}$, leaving slack within (84.1) for initialization, finite descriptor cases, and final wiring. Each lookup row needs at most eight comparisons, a fixed AND tree, a mask, and an accumulator addition. A required-site test uses two comparisons and a Boolean fold. A two-ended interval test uses two comparisons, one AND, and one addition. A compare-exchange uses a comparison and two signed selections. Every loop is over a fixed set of size $n, 2n, 3, \Delta, m, p, b$, or $2bc$. No loop is over F or coordinate extent. This allocation is not a claim that a macro has exactly its displayed number of gates.

84.1 Independent per macro construction check

The independent audit gives the following more concrete accounting, in the primitive units of Section 79. Equality costs at most four primitives; interval membership three; interval membership with a running count four; absolute value three; and an ordinary signed compare-exchange three. A fieldwise guard row needs three equalities, two conjunctions, a constant-bit mask and an accumulator fold, at most sixteen primitives.

For one of a descriptor's $2n$ raw slots, the audit allocates: one primitive for the key; $15n + 12$ for three presence scans and their masks; $4n + 7$ for raw-window count and size equality; $4n + 6$ for isolation; $20n$ for all competing raw slots; $22n + 19$ for the two context scans and their totalization; $32bc + 3$ for all guard rows and noncontextual bypass; and five for active/masking logic. These total at most $65n + 53 + 32bc$. Across $2n$ raw slots, then allowing $60n^2$ for all displacements and $2n^2 + 7n$ for outputs, sorting and the changing flag, one candidate-effect bound is

$$200n^2 + 120n + 64nbc + 20.$$

Multiplication by C remains below its $1000(C+1)(n+1)^2S_{\text{tab}}$ macro allocation. Retaining a separate branch field instead of a combined guard identifier replaces $64nbc$ by $84nbc$, still within the allocation.

Writing $P_2 = \binom{n}{2}$, a deliberately wasteful source/candidate/descriptor allowance is

$$P_2(200n + 200b + 100\Delta + 1000) + C(2000 + 200S_{\text{tab}}).$$

The second term permits up to 120 primitives for each of ten finite descriptor families, 200 for the two three-position sorters, 100 for partitions, 100 for fixed wiring, and $200S_{\text{tab}}$ for record selection. For $n \geq 2$, the inequalities $C \geq 3P_2$ and $C \geq \Delta P_2$ place this allowance below $1000(C+1)(n+1)S_{\text{tab}}$. If $n < 2$, these loops are empty. The stable minimum and state update need at most $C(n+12) + 10n + 40$ primitives, well inside their separate macro allocation. Initialization is $O(n)$ and final checks are in Q . These explicit loop charges support the conservative 10000 prefactor without pretending to have mechanically measured a general exporter.

Equations (79.4)–(79.5) imply

$$\begin{aligned} V &\leq 4U, & E &\leq 3U + Q, \\ L_1 &\leq 10U + 4Q, & L_2 &\leq 68U + 16Q. \end{aligned}$$

At most $Q \leq 3n + 1$ additional checks suffice for n external sign constraints, $n - 1$ strict-order checks, n target equations, and one completion equation. Their comparison gates are included in U . On an externally promised sorted canonical input domain, omit the admissibility checks. The $T = 0$ specialization is accounted separately.

84.2 Source preprocessing is separate

The pinned validator visits all bc branch/class pairs, evaluates guard syntax, and constructs the full domain and image tables plus incidence metadata. Retaining that implementation costs $O(cG_{\text{src}} + bc)$ scalar guard-program operations and a coarse $O(bc^2)$ bit-work bound for growing masks, in addition to parsing and integer-comparison costs. Source syntax and $2bc$ table entries occupy actual storage. Large binary J is not normalized cheaply; there is no polynomial-in-log J preprocessing claim.

84.3 The universal-source count is only a slot count

For the pinned universal source,

$$m = 122622, \quad b = 141561, \quad J = 0, \quad \Delta = 2, \quad F = 269291358255.$$

At $n = 5$, formula (80.1) gives

$$C = 10(2 + 3 + 12) = 170, \quad S_{\text{tab}} = 1396672.$$

This is *170 candidate slots*, not 170 witnesses, equations, or operations. A generic per-use multiplexer still pays the full source and guard-table costs. The resulting crude U is enormous; the theorem does not claim a practical universal quartic exporter. It removes the obligation to allocate one witness block for every factor position in every literal CA step, while preserving the costs actually used by this construction.

85 Coordinate heights and bit complexity

Let $H_0 = \max_i |x_i|$, with $H_0 = 0$ for the empty input. On an accepting execution, each matched particle displacement in one changing factor is at most $2B_3$. Actual snapshots therefore have coordinate magnitude at most $H_0 + 2KB_3$; hypothetical candidate outputs need one additional such displacement. Keys, shape offsets, and context endpoints add $O(B_3 + Z + J)$. Indices are at most F , time is at most T , and particle counts are at most $2n$.

For the specific totalized construction above, a coarse bound on the magnitude of every canonical auxiliary is

$$\text{Ht}_* = 100(n + 1)(H_0 + 2(K + 2)B_3 + Z + J + F + T + K + 1). \quad (85.1)$$

It covers comparison slacks, constant quotient/remainder, source codes, factor indices, class offsets, masked displacements, count accumulators, and canonical signed parts. Before dependent arithmetic, invalid modes and travel parameters are masked to fixed valid dummies. Guard lookup is fieldwise. A stable first-hit scan avoids accumulating many unbounded context offsets. Source-record selectors are one-hot, so partial selected-field sums are bounded by the largest selected field. Remaining multiplication is by a Boolean, sign, or the fixed factor-index constants; there is no hidden product of arbitrary coordinate magnitudes.

[write] Report 20 writes this bound H_* . It is printed Ht_* here because Report 28 (Section 119) uses H_* for an isolation distance; nothing else changed.

Thus $W = \lceil \log_2(\text{Ht}_* + 1) \rceil$ bits per auxiliary suffice, and the complete witness uses at most VW bits, beyond the external input encoding. This accepting-fiber bound is not a promise about adversarial submitted values. A rejected budget can still be evaluated using an analogous bound with R in place of K .

Let λ be the largest actual residual coefficient bit length. Trace-generation constants are at most a fixed multiple of $F + Z + J + T + 1$. Prescribed target constants appear only in acceptance equations and are charged separately in λ . Symbolic coordinate inputs do not make coefficients depend on H_0 ; specializing the input coordinates to constants instead must charge their bit lengths in λ .

The residual coefficient list has $O(U + Q)$ terms. With schoolbook arithmetic, straight-line witness generation costs

$$O((U + Q)W^2)$$

elementary bit operations, with source-constant lengths included in W , plus preprocessing. Direct residual and sum-of-squares verification is bounded by

$$O((U + Q)(W + \lambda + \log(U + Q + 2))^2).$$

For an arbitrary submitted tuple, use its actual maximum input/auxiliary bit length in place of W , or reject an oversized encoding after checking its length. This charges intermediate products, squaring, and accumulation rather than treating giant arithmetic as unit cost.

If every residual coefficient has magnitude at most C_0 , each collected expanded coefficient has magnitude at most $L_2 C_0^2$ and bit length at most

$$2\lambda + \lceil \log_2(L_2 + 1) \rceil + 1.$$

The full expansion enumerates at most L_2 ordered term occurrences and collects matching degree-at-most-four monomials. Its output storage and coefficient accumulation are separate costs; no equality of factored and expanded representation sizes is assumed.

86 Complete implemented mass two example

The source has controls **q**, **halt**, $J = 0$, and one branch from **q** to **halt** with side +1, update +1, and true guard. It has $D = 8$ and $F = 95$. Triple endpoints cannot occur at mass two. The only potentially changing factors are the outbound/inbound free E factors at indices 0, 22 and free P factors at indices 47, 70. The exporter specializes those four exact factor actions directly; it is not a general descriptor compiler.

For this implementation only, the symbolic input uses three natural coordinates

$$(h^+, h^-, g), \quad h = h^+ - h^-, \quad h^+ h^- = 0, \quad X = (h, h + g + 1).$$

The gap-minus-one encoding enforces strict order without an input comparison. It differs from the general theorem's four-coordinate canonical-pair interface but represents every sorted mass-two support uniquely. All endpoint values are checked to be exact integers and strictly increasing before they become coefficients; Boolean and floating-point endpoints are rejected.

The normative sample uses $X = (-10, -5)$, $T = 1$, $K = 2$, and target $Y = (-9, -4)$. Its input values are $(0, 10, 4)$. The deterministic execution is

Round	Action	Support afterward	Time	Cursor
0	Factor 0	$(-9, -3)$	0	1
1	Factor 47	$(-9, -4)$	0	48
2	Absence and completion	$(-9, -4)$	1	0

The third round is necessary even though both permitted changes have already occurred. A budget $K = 1$ rejects this input for one full CA step.

The emitted primitive ledger and polynomial counts are

Quantity	Exact sample value
Arithmetic assignments A	189
Comparisons I	78
Constant divisions D_v	0
Boolean gates G	78
Acceptance/admissibility residuals Q	4
Natural auxiliary variables V	612
Quadratic-or-lower residuals E	616
Residual monomials	2013
Ordered expanded-square term occurrences	9955
Collected expanded quartic monomials	4573
Largest witness bit length	7
Largest residual coefficient bit length	7

The residuals include external sign separation, horizon completion, and two endpoint equations. Formula (79.4) gives $378 + 156 + 78 = 612$ and (79.5) gives $378 + 156 + 78 + 4 = 616$. These values describe this particular small source and sample, not the general schema or universal-source instantiation.

86.1 A coefficient explicit artifact rather than an abbreviation

The bundled `example-polynomial.json` contains every residual coefficient, every variable name in declared order, and the entire collected expanded quartic. A term is serialized as

$$[\text{integer coefficient}, [i_1, \dots, i_d]],$$

meaning the coefficient times $v_{i_1} \cdots v_{i_d}$; an empty index list denotes a constant. Indices are sorted, zero coefficients omitted, and repeated monomials collected. The first three variables are external inputs and the next 612 are the auxiliaries. The companion `example-witness.json` gives all 612 natural values. These complete machine-readable files, rather than a many-page typesetting of 4573 terms, are the normative coefficient and witness listing.

The binding verifier regenerates the normative small circuit and checks the source, inputs, T , K , target, arity, variable order, ledger, residual rows, expansion, round decoder, and output. It rejects malformed types and noncanonical encodings. It then evaluates the supplied witness exactly. This is binding verification using the producer's small compiler; the separate audit independently reconstructs and checks the expanded polynomial and its natural witness. Neither claims an independent arbitrary-mass exporter.

87 Reproducibility and audit evidence

The release contains the complete finite-schema proof, fixed-slot scheduler, mass-two emitter, coefficient and witness files, exact receipts, independent audit, pinned evaluator and compiler, the compressed universal source needed by the count check, and dependency proof references. Python replay uses only the standard library. It does not fetch packages or resolve optional private files.

[write] The compressed universal source `reproducibility/source.json.gz` is not shipped in this directory, and neither are the bundled copies of Report 19’s evaluator and Report 15’s compiler (shipped once, as `code/19-lazy-eval-lazy_reversible.py` and `code/15-rev-five-compiler-reversible_binary.py`) nor the dependency references; the README says how to rebuild or retrieve them (“Reconstructing the excluded data”).

The general scheduler tests compare candidate sets and results with the pinned lazy/eager factor products on small sources. Cases include empty and singleton supports, simultaneous separated copies, the known four-event malformed cascade, exact and insufficient budgets, exact completion-round counts, canonical idle padding, and multistep total budgets. The mass-two suite compares emitted polynomials with actual CA iteration, including 100-digit signed coordinates, and checks the complete sample expansion and every one-coordinate +1 witness mutation. Guard checks cover actual finite tables and boundary conditions. These are finite executable evidence, not a substitute for the general theorem.

The independent audit implements the full $2n$ -slot raw/absence/context/displacement algorithm separately and compares it with literal original factor application. Its seven sources include no branches, direct-only behavior, left decrement, guarded increment, mixed original branch order, $J = 17$ context, and multiple incidence branches. It checks 942 configurations, 123222 factor comparisons, 468 scheduler cases, and 71 strict underbudget rejections. It exercises 485 multi-active factor instances, 20 competitor exclusions, and 14 multiple-particle context rejections. Arithmetic checks include 392 exhaustive signed-division fibers, 36 sample syntax/ledger cases, and 189 simultaneous positive/negative-part mutations. Final interface replays reject 13 malformed targets, five other malformed inputs, and ten malformed serialized/source cases; a valid 301-digit-coordinate trace passes.

The packaged replay runs the applicable suites in ordinary and optimized Python modes; no mathematical rejection or validation depends on `assert`. File binding and exact integer checks are explicit. Conservative resource and integer-encoding limits keep this a practical sample-release checker rather than an unrestricted service for adversarial polynomials. Finite tests corroborate the construction but do not prove universality or replace its mathematical argument.

87.1 Pinned dependencies and provenance

The evaluator’s byte identity is

42e8aa65c05fcf373a03a02be51ebb89a4fdcb1f1070ad776ffe1e8e99049c61

The original frozen compiler’s byte identity is

f95a6028c9d3cc6f8b0b497cfc0cad68104da6d70e55205df04ad234e5d0c76f

The package verifies these mandatory local dependencies before execution. Any portability edits to research drivers are recorded as original-versus-adapted hashes; the frozen rule implementation is not silently modified. Exact file inventories and byte lengths are checked, with extra-file and symlink rejection. A manifest is an integrity record tied to the release; it is not a cryptographic author identity or an external proof of authenticity.

The README gives the authoritative replay and build commands, expected outcomes, implementation limits, and manifest policy. An extracted copy is replayed independently of the authoring directories before delivery. LaTeX source is supplied for the report; rebuilding the PDF needs TeX Live, while numerical replay does not. Report 8’s canonical comparison and signed-division constructions are included as local references [18]; no MRDP existence theorem is invoked.

88 Limits and the next concrete improvement

A bare list of changing factors proves too little. It does not certify that an earlier omitted factor was an identity, that a remaining suffix has no change, or that all active occurrences of the selected factor were applied simultaneously. The present construction pays for those obligations by evaluating the whole fixed candidate circuit from each snapshot.

The cost is event-budget-sensitive but is not merely proportional to active keys or to the operational lazy evaluator's number of tested factors. There are $T + K$ complete candidate rounds, including losing candidates and ghost computations in padded rounds. To inherit a smaller data-dependent lazy-test count, one would need a different explicit sequential-search transcript and its own fully paid bound.

No unconditional small numerical K bound follows for arbitrary malformed finite supports: cascading changes are possible. F remains present through factor-index lengths and source-dependent coefficients, though not as an allocated factor array. The universal mass-five candidate count does not show universality for mass two or a practical universal polynomial exporter. There is no claim of fixed quantified arity for unbounded horizons, of a finite-fold representation of all recursively enumerable sets, of a real-witness equivalence, or of a new optimality or novelty result.

[write] (3 October 2026.) This step is still open for the ordered rule. Report 28 (Part [VI](#)) ships a general emitter with measured ledgers, but for the new parallel rule of Report 26, which differs from the ordered rule on malformed supports (Section [120.3](#)); it is not an emitter for Theorem [77.1](#).

The next engineering step is a general arbitrary-mass coefficient emitter implementing every explicit macro, with an independent binding checker for those macros and measured rather than conservative counts. The present release establishes the finite canonical-witness theorem and provides a complete smaller instance that can already be expanded and checked exactly.

Part VI

The parallel two-involution rule: replacement, sparse evaluation and quartic certificates (Reports 26, 27 and 28)

[write] The three Reports of this Part concern one new rule. Report 26 replaces the ordered product of Report 15’s compiler by two parallel full-shift involutions, $F_{\text{new}} = P_{\parallel} \circ E_{\parallel}$, that agree with the old rule on every admissible encoded state and deliberately differ on malformed ones; Report 27 evaluates the new rule exactly on every finite support without building its endpoint family; Report 28 compiles it, at each fixed horizon, into a degree-four polynomial with a unique natural witness. The three Reports write the same map as F_{new} , Φ and $\mathcal{F}$ respectively. The Hilbert’s-tenth programme reviewed all three in `Computability/HilbertTenthProblem/Papers/research-wip/native-stream-queue/review_parallel_particle_reports.md` (commit `fb7e3cb47`): it read their main proofs, enumerated all 4,096 period-12 words of Report 26’s lemma example, reconstructed Report 28’s 12,595-term quartic from its 1,502 residuals, found no defect, and records that the horizon-dependent arity gives no fixed-arity, finite-fold or universal result.

Two parallel conservative involutions

A full shift replacement for the five particle compiler

Research report 26

3 October 2026

Abstract of Report 26, as delivered. We replace the ordered local-gate construction of a binary five-particle cellular automaton by two parallel involutions. A general prospective-isolation lemma gives a full-shift, number-conserving involution of radius at most $3(b+r)$ from finitely many endpoint rewrites of write radius b and candidate-read radius r . The selection test preserves the complete set of raw type-anchor keys, and a read/support separation argument proves invariance of every eligibility status, including initially rejected keys. Applied to the literal five-particle templates, the two blocks agree with the previous construction on its entire admissible doubled micrograph. The encoding, exact clocks, reflections, halt observer, and clean targets are unchanged. An explicit sufficient composite radius is $180D + 258 + 9J$, with symbolic universal-ledger value 91 711 698. The resulting rule intentionally differs on malformed inputs. We include the boundary proof, finite executable evidence, precise attribution, and a portable sealed replay package. We claim neither radius optimality nor novelty of the two-involution architecture; no universal new-rule evaluation or truth table is produced.

[write] Delivered as `Two_Parallel_Conservative_Involutions_Package.zip` (batch 82, archive 20; arrival `db37d18c8`), dated 3 October 2026; files placed by `7d2b1b245` with the prefix `26-parallel-involutions-`. Report 26 keeps the endpoint templates of Report 15’s compiler (Part I) and replaces their ordered execution; its universal ledger is Report 19’s (Part IV). Lemma 90.1 is a lemma proved here, in the marker and conserved-landscape tradition that Section 98 names; Report 26 disclaims priority for it, and so does this report. Report 27 restates it (Section 101ff.). Report 26’s macro `\N` (printing $\mathbb{Z}_{\geq 0}$) is `\Nge` here; nothing printed changed. The delivered table of contents is omitted.

89 Main result and the object that changes

A configuration is a word in the full binary shift $\{0, 1\}^{\mathbb{Z}}$. Its particles are its sites with value one. A cellular automaton (CA) here is a translation-equivariant map whose output at each site depends on a uniformly bounded neighborhood. A CA is number-conserving if it preserves the number of particles on every finite-particle configuration. No assertion equating divergent infinite sums is intended.

The starting point is the finite guarded-swap compiler of Report 15, pinned in the accompanying package as `scientific/frozen_reversible_binary.py`. Its SHA-256 is

f95a6028c9d3cc6f8b0b497cfc0cad68104da6d70e55205df04ad234e5d0c76f.

That compiler produces an ordered list of individually reversible gates. Its old edge and phase blocks are products of many such gates. They are matchings on admissible encodings, but need not be involutions on malformed configurations. We keep its finite endpoint-template data and replace its global selection and execution mechanism.

Theorem 89.1 (Source-uniform replacement). *For every finite source accepted by the pinned compiler, there is a binary CA*

$$F_{\text{new}} = P_{\parallel} \circ E_{\parallel}$$

with the following properties.

1. *Each block is an involution on all of $\{0, 1\}^{\mathbb{Z}}$, is translation-equivariant, and conserves ordinary finite particle number.*
2. *The new rule agrees step for step with F_{old} on the entire admissible doubled micrograph, including every natural-counter home, both signs of every strict intermediate node of every enabled edge, and all translations.*
3. *It has sufficient radius*

$$R_{\text{new}} = 180D + 258 + 9J, \quad D = 2m + 4p,$$

where m is the number of source controls, p the number of nonzero-update branches, and J the finite counter-class cut.

4. *Its inverse is $E_{\parallel} \circ P_{\parallel}$. It is time-symmetric, with either of the individually number-conserving binary involutions $E_{\parallel}, P_{\parallel}$ as a time reverser.*

The proof occupies Sections 90–93; the radius calculation is in Section 95. Equality on admissible states is a simulation theorem, not equality on the full shift. Section 96 gives an explicit malformed input on which the old and new rules differ. Thus the displayed improvement concerns a *new extension of the same encoded dynamics*. It is not a smaller-radius presentation of the previous full-shift automaton.

The statements about a universal source are inherited conditionally through the previously established source construction and its admissible encoding. This report proves the replacement uniformly for all accepted sources; it does not re-prove the attributed universal-source theorem, materialize a universal template array, or execute that array. Mathematical existence, finite examples, and practical rule evaluation are kept distinct throughout.

90 Prospective isolation on the full shift

90.1 Write supports and read windows

For $a \geq 0$ and $u \in \mathbb{Z}$, put $I_a(u) = [u - a, u + a] \cap \mathbb{Z}$. Fix integers $0 \leq b \leq r$ and a finite type set G . Each type g has its own write set

$$W_g \subseteq I_b(0)$$

and two distinct binary words p_g, q_g on W_g of equal Hamming weight. Define $\tau_{g,u}$ everywhere on $\{0, 1\}^{\mathbb{Z}}$: exchange p_g and q_g on $u + W_g$, leave every other word on that write set fixed, and never change coordinates outside $u + W_g$. Then $\tau_{g,u}^2 = \text{id}$.

The common bound b is *not* a requirement to erase or constrain every bit of $I_b(u)$. Bits in $I_b(u) \setminus (u + W_g)$ are retained. This distinction matters below because a free pair has a smaller exactness window than the common triple write bound.

Let $c_g(x, u)$ be a translation-covariant Boolean predicate depending only on $x|_{I_r(u)}$. Assume that candidacy implies that the word on $u + W_g$ is p_g or q_g , and that the predicate is endpoint-symmetric:

$$c_g(x, u) = c_g(\tau_{g,u}x, u). \quad (90.1)$$

Extra empty-padding tests or contextual guards are permitted, provided all their reads are included in r . A *raw key* is (g, u) ; define

$$C(x) = \{(g, u) : c_g(x, u) = 1\}.$$

Different types at the same anchor are different keys. Endpoint orientation is not part of the key.

90.2 Selection without a recursive eligibility test

Set $H = 2(b + r)$. A raw key $k = (g, u)$ is selected precisely when

$$k \in C(x), \quad (90.2)$$

$$\text{there is no } (h, v) \in C(x) \setminus \{k\} \text{ with } |v - u| \leq H, \quad (90.3)$$

$$C(\tau_k x) \cap (G \times I_{b+r}(u)) = C(x) \cap (G \times I_{b+r}(u)). \quad (90.4)$$

The last test compares *all raw types*, before and after this one hypothetical rewrite. It does not compare only the current type, only the currently selected keys, or oriented endpoint labels. It does not recursively ask whether other keys are eligible.

Write $A(x)$ for the selected-key set, and let $T(x)$ apply the selected rewrites simultaneously. Distinct selected anchors have distance greater than $H \geq 2b$, so their containing write intervals are disjoint. Consequently this definition makes sense even on an arbitrary bi-infinite dense input: an output site participates in at most one selected rewrite.

Lemma 90.1 (Prospective-isolation lemma). *For every $x \in \{0, 1\}^{\mathbb{Z}}$,*

$$C(Tx) = C(x), \quad A(Tx) = A(x), \quad T^2x = x.$$

The map T is a translation-equivariant, number-conserving CA of radius at most $3(b + r)$. The type count does not enter this radius bound.

Proof. We give all eligibility-invariance steps, since preservation of already selected keys alone would not prove involutivity.

A finite prospective check is global raw-key preservation. A rewrite at u can affect a raw predicate at v only if its write set meets $I_r(v)$. Necessarily $|v - u| \leq b + r$. Thus (90.4) is equivalent to

$$C(\tau_k x) = C(x). \quad (90.5)$$

This equivalence concerns raw predicates, not eligibility.

No cooperative births occur. If one raw read window $I_r(v)$ met the supports of two selected rewrites at u, w , then

$$|u - w| \leq |u - v| + |v - w| \leq 2(b + r) = H,$$

contrary to isolation. Every raw predicate after the simultaneous update therefore sees either its original word, or exactly the word obtained after one selected rewrite alone. Equation (90.5) preserves its value in either case. This proves $C(Tx) = C(x)$ for every raw key, including keys initially absent.

The entire prospective read neighborhood is controlled. Write $P_k(x)$ for the equality in (90.4). Candidate anchors examined there lie within $b + r$ of u ; each candidate reads radius r . The whole test is determined by

$$x|_{I_{b+2r}(u)}. \quad (90.6)$$

The hypothetical rewrite needs no coordinates outside this interval. Since $\tau_k^2 = \text{id}$, interchanging the two sides of the comparison gives

$$P_k(\tau_k x) = P_k(x). \quad (90.7)$$

Every eligibility status is invariant. Raw-key invariance already preserves isolation. Consider any isolated raw key $k = (g, u)$, whether selected or not. Every other selected key (h, v) is a different raw key, so $|v - u| > 2b + 2r$. Its write interval is disjoint from (90.6). If k was rejected, its full prospective-test neighborhood is therefore unchanged. If k was selected, only its own rewrite affects that neighborhood, and (90.7) preserves the test. Nonisolated raw keys remain nonisolated; absent keys remain absent. This proves $A(Tx) = A(x)$, including the initially ineligible cases.

The second update is the inverse. The second application uses exactly the same disjoint type-anchor rewrites. Each is an involution, so every changed block returns to its original word and every other site remains fixed.

Locality and conservation. At anchor u , candidacy reads radius r , isolation reads radius $H + r$, and prospectivity reads radius $b + 2r$. Their maximum is $2b + 3r$. Only keys within distance b can write a given output coordinate. Hence

$$R_T \leq b + \max\{r, H + r, b + 2r\} = 3(b + r).$$

There are finitely many types and finitely many anchors in every query, so this is a finite Boolean local rule. Translation covariance gives shift equivariance. Disjoint equal-weight rewrites conserve particle number on every finite input. $\square$

For infinite inputs the same proof gives a well-defined local bijection. One may match input and output particles within each changed finite block to obtain bounded-displacement local transport; finite total mass is not required for the involution proof. The independent full-shift audit is frozen as `scientific/audit-lemma.md`.

[write] The audit is shipped as `26-parallel-involutions-audit-lemma.md`. The hypothesis $b \leq r$ of Section 90 is not used in the proof (observed at placement; Report 27 states the lemma with $r \geq b$, Report 28 without the order). The Hilbert's-tenth review `review_parallel_particle_reports.md` (fb7e3cb47) re-derives the lemma and the radius $3(b + r)$, and checks the period-12 example of Section 96 on all 4,096 words.

[write] Added 4 October 2026, batch 91 (reciprocal note): Report 49 of the same pipeline, printed as source 24 in Part IX of `signal-machine-collision-certificates` (written in commit `61c9e3eb2`), proves this lemma again as its Lemma `smc:rt:lem:guard`, with the same isolation distance $H = 2(b + r)$ and the same radius bound $3(b + r)$, for its own row format, in which a key does not record its endpoint side. It credits the mechanism to the review named above, and uses it for a globally reversible binary rule, of forward and inverse radius at most 90, that realizes the four-particle timing threshold. That report prints it as a marked second route, with no novelty claimed.

91 The finite source and five particle geometry

91.1 Source interface and admissibility

An accepted source is a finite partial-injective deterministic machine on controls Q and counters $(c_0, c_1) \in \mathbb{Z}_{\geq 0}^2$. A branch $e : q \rightarrow q'$ selects the left or right counter, represented by $v \in \{-1, +1\}$, and changes it by $\Delta \in \{-1, 0, +1\}$. Its Boolean guard $G_e(c)$ uses finitely many tests $c_j = k$ or $c_j > k$. Enabled updates must preserve natural counters. Domains of branches from each control are disjoint, and *actual images* of branches entering each control are disjoint. The designated halt control has no outgoing branch.

For selected counter index j , the exact image predicate is

$$I_e(c') = G_e(c' - \Delta \mathbf{e}_j) \wedge (c' - \Delta \mathbf{e}_j \in \mathbb{Z}_{\geq 0}^2).$$

The cut J ensures that all domains and images are determined by the two counter classes $0, 1, \dots, J, > J$. The pinned parser validates this finite interface, including unreachable rows; it does not accept a caller-supplied unchecked image predicate. Its complete schema is supplied as an inert reference.

There are $m = |Q|$ home modes H_q and two moving modes O_e, I_e for each of the p nonzero branches. Give every signed mode a distinct gap $d(u^+), d(u^-) \in \{1, \dots, D\}$, where

$$\begin{aligned} D &= 2m + 4p, & S &= 2D + 2, & B_2 &= D + 1, & L &= 3D + 4, \\ b &= B_3 = 4D + 5 = L + D + 1, & K &= 3b + 1, & Z &= 10b + 10 + 2J. \end{aligned} \tag{91.1}$$

A head at anchor x is the pair $\{x, x + d(u^\pm)\}$. Its anchor is always the leftmost occupied site, including during leftward motion. With origin zero, a home encoding is

$$\text{Enc}(q, c_0, c_1)^\pm = \{-(Z + c_0), 0, Z + c_1, S, S + d(H_q^\pm)\}. \tag{91.2}$$

There are exactly five particles, no extra track, no spatial parity background, no clock particle, and no mass-free phase symbol.

91.2 The unsigned micrograph

A zero-update source edge is one home-to-home microedge. A nonzero enabled edge is subdivided as follows, with origin o , selected old marker m , and new marker $m' = m + v\Delta$:

1. Dispatch H_q at $o + S$ to O_e at $o + vS$, guarded by G_e .
2. Move the outbound anchor by v until it reaches $m - vS$.
3. Move the selected marker from m to m' and change the head to I_e at $m' - vS$.
4. Move the inbound anchor by $-v$ until it reaches $o + vS$.

5. Commit to $H_{q'}$ at $o + S$, guarded by I_e on the post counters.

The admissible unsigned graph contains every home ID over natural counters and only the strict intermediate configurations of these subdivisions of *enabled* source edges. Geometrically plausible O/I packets of a disabled branch are excluded. The admissible doubled graph contains both signs of every such node and every translation.

The branch label and the old counters in O_e , or post counters in I_e , determine the source edge; the head anchor determines the position on its corridor. Thus distinct edges cannot share an intermediate ID. Home merging is injective because the actual branch images are disjoint. This is a partial-injective micrograph. Its doubled edge matching exchanges c^+ with c'^- for each unsigned edge $c \rightarrow c'$, and the phase matching exchanges the two signs at each node.

92 Literal templates and the new two blocks

We now give the full finite family, so the construction does not depend on an unspecified particle mechanism. Every displayed pair P_g, Q_g is a pair of occupied supports, understood as binary words on the type-specific interval $W_g = I_{B_g}(0)$. Pair types have $B_g = B_2$; triple types have $B_g = b$. The invariant template anchor is zero in the display. It need not be the head's leftmost site or an occupied site in both endpoints.

92.1 Edge templates

For each moving mode u , let w be its forward direction (v for O_e , $-v$ for I_e). There is one free pair type:

$$\{0, d(u^+)\} \longleftrightarrow \{w, w + d(u^-)\}. \quad (92.1)$$

Its invariant key is the *predecessor* head anchor. Near a marker, there are triple types

$$\{0, wt, wt + d(u^+)\} \longleftrightarrow \{0, w(t+1), w(t+1) + d(u^-)\}, \quad S \leq t \leq L, \quad (92.2)$$

$$\{0, -wt, -wt + d(u^+)\} \longleftrightarrow \{0, -w(t-1), -w(t-1) + d(u^-)\}, \quad S+1 \leq t \leq L. \quad (92.3)$$

For each nonzero branch there are dispatch, endpoint, and commit triples:

$$\{0, S, S + d(H_q^+)\} \longleftrightarrow \{0, vS, vS + d(O_e^-)\}, \quad (92.4)$$

$$\{0, -vS, -vS + d(O_e^+)\} \longleftrightarrow \{v\Delta, v\Delta - vS, v\Delta - vS + d(I_e^-)\}, \quad (92.5)$$

$$\{0, vS, vS + d(I_e^+)\} \longleftrightarrow \{0, S, S + d(H_{q'}^-)\}. \quad (92.6)$$

Dispatch uses G_e , endpoint has no contextual guard, and commit uses the exact I_e . The endpoint key remains the *old* marker coordinate, even after the marker moves. Each zero-update branch contributes

$$\{0, S, S + d(H_q^+)\} \longleftrightarrow \{0, S, S + d(H_{q'}^-)\}, \quad (92.7)$$

guarded by $G_e = I_e$.

92.2 Phase templates

For each moving mode, use the pair

$$\{0, d(u^+)\} \longleftrightarrow \{0, d(u^-)\},$$

and, for $s \in \{-1, +1\}$ and $S \leq t \leq L$, the triples

$$\{0, st, st + d(u^+)\} \longleftrightarrow \{0, st, st + d(u^-)\}.$$

For every control use the home triple

$$\{0, S, S + d(H_q^+)\} \longleftrightarrow \{0, S, S + d(H_q^-)\}.$$

All phase types are unguarded.

92.3 Raw predicates and guard reads

For a type with write radius B_g , put $\ell_g = 3B_g + 1$. Its new raw predicate at u requires that all particles in $I_{\ell_g}(u)$ be exactly $u + P_g$ or $u + Q_g$, and that its old contextual guard, if any, hold. We omit the old *same-type raw-key exclusion*. The new all-type isolation and prospective test replace that safety mechanism; the template list is no longer an execution order.

At home interactions the counter classes are read at

$$u - (Z + k), \quad u + (Z + k), \quad 0 \leq k \leq J.$$

One occupied position selects a finite class, all-zero reads mean $> J$, and multiple occupied positions make the guard false. These reads lie outside the type's write support. The endpoint swap leaves them untouched, and it preserves exactness in $I_{\ell_g}(u)$, so every raw predicate is endpoint-symmetric.

For the edge family take $r_E = Z + J$; this dominates every exactness and guard read. For the phase family take $r_P = 3b + 1 = K$. Lemma 90.1 gives the two full-shift involutions $E_{||}, P_{||}$, each using the common write bound b . Empty edge families give the identity block. Sources with no nonzero branch are covered without adding fictitious moving modes.

93 Preservation on every admissible state

93.1 Unique recognition

Every admissible state has one unique pair of particles at distance at most D : its head. Markers have mutual separation at least Z ; head-to-marker distances are at least $S - D = D + 2 > D$. Every endpoint template has exactly one close pair and no other distance at most D . Therefore a raw match must use the actual head pair. Its globally unique gap fixes its signed mode, and a triple match fixes the relevant marker. No background marker pair or head-marker pair can replace the head.

The diameter of a triple endpoint is at most $2b < Z$, so two different markers cannot both serve as its nearby singleton. For each individual old type there is at most one old raw key on an admissible state. Thus its old same-type exclusion was vacuous there: removing that test introduces no hidden matches on valid states.

The other markers lie outside the intended triple's exactness radius K . Ordinary triple keys are markers, at distance at least $Z > K$ from other markers. At a reverse endpoint interaction the key is the old marker location, at distance at most one from the new marker; the other markers are still at least $Z - 1 > K$ away. Hence the intended triple exactness tests hold.

93.2 The complete corridor boundary partition

For a moving mode, orient its corridor in direction w . Let A be the departure marker, C the destination marker, and let x be the head anchor. Set

$$N = w(C - A) \geq Z, \quad t = w(x - A), \quad S \leq t \leq N - S.$$

For O_e the corridor length is $Z + c$; for I_e it is $Z + c + \Delta$. The third marker is outside every relevant near-contact range, and $N > 2(L + 1)$.

For a plus state the free key is $u = x$. Pair exactness holds exactly when $L < t < N - L$. The plus cases are the disjoint exhaustive partition

Candidate	Oriented head position
Behind triple	$S \leq t \leq L$
Free pair	$L + 1 \leq t \leq N - L - 1$
Ahead triple	$N - L \leq t \leq N - S - 1$
Forward interaction	$t = N - S$

For a minus state the free inverse key is $u = x - w$, whose oriented coordinate is $t - 1$. Pair exactness holds exactly when $L < t - 1 < N - L$. The minus partition is

Candidate	Oriented head position
Reverse interaction	$t = S$
Behind inverse triple	$S + 1 \leq t \leq L + 1$
Free inverse pair	$L + 2 \leq t \leq N - L$
Ahead inverse triple	$N - L + 1 \leq t \leq N - S$

Both tables include the equality cases. A behind triple from distance L reaches minus distance $L + 1$, but the free predecessor key still sees its marker at distance L , so a free inverse cannot compete. An ahead free move from marker distance $L + 1$ reaches minus distance L , but there is no ahead triple with predecessor distance $L + 1$. At forward arrival distance S there is no further travel, and no travel from behind distance $S - 1$ enters a reverse interaction. These facts hold for both physical directions; the head anchor convention has not been reflected.

93.3 Every interaction and every guard

At home plus, only outgoing dispatch/direct types of that control can match. Their domain predicates are disjoint and test the actual counters, so there is one candidate exactly when the source ID has a successor. At home minus, only incoming commit/direct reverse endpoints can match. Their *image* predicates are disjoint and include natural pre-counter existence, so there is one candidate exactly when that ID has a predecessor. Using a forward predicate at a reverse commit would invalidate this conclusion.

At the outbound forward wall, the endpoint rewrite is anchored at old marker m , and produces marker $m' = m + v\Delta$ and inbound-minus head anchor $m' - vS$. The reverse endpoint is necessarily anchored at

$$m' - v\Delta = m,$$

because its singleton offset is $v\Delta$ and its head offset is $v\Delta - vS$. The same type-anchor key survives the swap, even when the actual marker moves.

At the outbound reverse wall, the old counters satisfy G_e because this intermediate belongs to an enabled edge; reverse dispatch returns to its genuine home-plus predecessor. At the inbound forward wall, the post counters satisfy I_e for the same reason; commit reaches its genuine home-minus target. At the inbound reverse wall, the endpoint inverse restores the old marker and outbound-plus head. Branch-specific signed gaps eliminate all other interaction types, and the corridor partitions eliminate travel competition.

It follows that every admissible edge-block input has zero or one raw key globally. A present swap remains admissible and leaves exactly the same raw type-anchor key, with its orientation reversed.

A missing incident microedge leaves no raw key. This argument covers arbitrary counters and all enabled branches, not just a sampled forward orbit.

93.4 The phase block and agreement theorem

At a home, its unique phase triple flips sign. At a moving node, if every marker is farther than L from the unchanged head anchor, the phase pair is the unique match. Otherwise there is one marker at distance $S, \dots, L$, and its unique phase triple is the match. The marker and anchor geometry do not change during a phase flip, so exactly the same key remains, with opposite orientation.

A singleton raw-key set automatically passes all-type isolation. Since the corresponding swap preserves that singleton set, it also passes prospectivity. Empty edge sets remain empty. Thus $E_{\parallel}$ equals the old edge matching and $P_{\parallel}$ equals the old sign flip on every admissible signed node. Both maps preserve admissibility. Their composition and inverse therefore agree with the old rule and inverse throughout that set, proving the simulation part of Theorem 89.1. The source-uniform geometric audit is supplied as `scientific/audit-preservation.md`.

94 Clocks observations and time symmetry

94.1 Time reversal and the exact path

Since both new blocks are full-shift involutions,

$$F_{\text{new}}^{-1} = E_{\parallel} P_{\parallel}, \quad E_{\parallel} F_{\text{new}} E_{\parallel} = F_{\text{new}}^{-1}, \quad P_{\parallel} F_{\text{new}} P_{\parallel} = F_{\text{new}}^{-1}.$$

Each displayed time reverser is itself binary and number-conserving. This statement is immediate algebraically from the proved factorization; it does not attribute novelty to that algebraic architecture.

On an admissible plus node with successor, F_{new} advances one unsigned microedge and retains plus. On a minus node with predecessor it moves backward one microedge and retains minus. At a missing edge it changes sign in one step. There is no extra external phase schedule: the composition defines one time-independent CA step.

For a nonzero branch at selected old counter c , the two travel lengths are $Z + c - 2S$ and $Z + c + \Delta - 2S$. Adding dispatch, endpoint, and commit gives

$$\tau_e(c) = 3 + 2(Z + c) + \Delta - 4S. \quad (94.1)$$

A zero-update edge takes one step. Equality of admissible trajectories therefore preserves the exact old clocks, not only reachability or an asymptotic slowdown.

[write] Equation (94.1) is Report 15's clock (5); it is unchanged because the two rules agree on admissible states.

If the initial ID is predecessor-free and its forward path has T microedges, its reflected orbit has exactly $2T + 2$ states. Partial injectivity prevents an earlier repetition from a predecessor-free start. At halt the state reflects; it does not freeze or enter a separate attractor. Nonfinal stuck IDs also reflect. Consequently periodicity is equivalent to finite forward path length, and becomes equivalent to designated halting only with the additional premise that every nonhalting input has an infinite source path.

94.2 Halt observer and clean exact targets

The anchored halt word on $[0, S + D]$ has ones exactly at $0, S, S + d(H_h^+)$ and zeros elsewhere. Its length remains

$$S + D + 1 = 3D + 3.$$

On a valid orbit its occurrence is equivalent to forward reachability of the halt control. No backward traversal can introduce a new halt home, since that would require an outgoing edge from a halt control that has no exits.

For a clean whole-configuration target, assume the original source has a control-wide no-incoming start q_0 and a no-outgoing halt h . Use forward controls $F(q)$, backward controls $B(q)$, and a new terminal H . Include each original branch forward and its inverse backward with exact image guard, plus the zero-update bridges $F(h) \rightarrow B(h)$ and $B(q_0) \rightarrow H$. Domains and images exchange on reverse branches; the two empty boundary slots make the bridges compatible with partial injectivity. The wrapped source has

$$m' = 2m + 1, \quad p' = 2p, \quad a' = 2a + 2,$$

and the same cut J . All geometry constants must be recomputed for this enlarged source.

On input (a_0, b_0) the target is $\text{Enc}(H, a_0, b_0)^+$. It is computable and input-dependent, not a universal constant configuration. If Θ is the original forward physical microedge count with the *enlarged* geometry, then

$$\tau(c, \Delta) = \tau(c + \Delta, -\Delta)$$

and the first clean-target time is $2\Theta + 2$. The full reflected period is $4\Theta + 6$. For the bundled one-increment clean-target example, $D = 18$, $S = 38$, $Z = 782$, $\Theta = 1416$, first target time 2834, and reflected period 5670. The new radius is 3507; the five-particle orbit itself is unchanged.

The classical reversible-counter sources and their normalization belong to the inherited compiler argument, not a new source theorem in this report. Its proof and schema are included as references so these dependency boundaries remain auditable.

95 Radius and finite resources

Applying Lemma 90.1 independently gives

$$\begin{aligned} R_E &= 3(b + Z + J) = 33b + 30 + 9J, \\ R_P &= 3(b + 3b + 1) = 12b + 3. \end{aligned}$$

A composition has radius bounded by the sum, hence

$$R_{\text{new}} \leq 45b + 33 + 9J = 180D + 258 + 9J. \quad (95.1)$$

These are sufficient radii. The new bound is not pointwise smaller for every tiny source: for a two-control, one zero-update branch example with $D = 4$ and $J = 0$, the old bound is 540 and the new bound is 978. The improvement claimed for the universal ledger is a comparison of sufficient bounds for that source. No proof that either bound is minimal is offered. A sharper special-purpose separation argument could conceivably improve them.

[write] (4 October 2026.) For this source Report 70's bound (Theorem 127.1) is $108 \cdot 4 + 149 = 581$, still above 540, and Report 71's (Theorem 140.1) is $76 \cdot 4 + 105 = 409$, or 408 by its same-rule refinement, below 540. Each is a sufficient radius of a different full-shift rule with the same admissible dynamics; the sentence above remains true of Report 26's bound.

The literal family still has

$$8pD + 29p + m + a$$

endpoint types, where a is the number of zero-update branches. In the old construction this counted ordered factors: $4p$ pair gates, $8pD + 23p + m$ unguarded triples, and $2p + a$ contextual triples. Here the same data count template alternatives inside *two* parallel involutions. The inherited ledger fields named **E**, **P**, and **factors** retain type counts, not a new ordered runtime sequence.

The old sufficient radius was

$$R_{\text{old}} = 4p(6D + 8) + (8pD + 23p + m)(24D + 32) + (2p + a)(Z + J + 12D + 16). \quad (95.2)$$

The new bound depends on geometric neighborhoods rather than summing a radius once per executed template. This does not remove template dependence from the Boolean computation of candidacy or prospectivity.

Source or substitution	Old bound	New bound
Two-control increment, $D = 8, J = 0$	21 144	1 698
Clean-target sample, $D = 18, J = 1$	164 314	3 507
Universal ledger, $D = 509\,508, J = 0$	3 292 955 588 459 274 804	91 711 698

The universal row substitutes the frozen Report 19 ledger

$$m = 122\,622, \quad p = 66\,066, \quad a = 75\,495, \quad J = 0$$

into source-uniform formulas. The original ledger and its provenance are supplied as inert evidence. It has 269 291 358 255 endpoint types, and that array is *not* materialized here. The symbolic full local truth table would have 2^{2R+1} entries at a chosen sufficient radius, but no such table is emitted. The reference interpreter constructs small-source arrays for transparency and does not provide a universal-efficient new-rule evaluator.

96 Malformed inputs and why prospectivity matters

For controls q, h , a true-guard right increment, and $J = 0$, consider

$$X = \{-118, -112, 0, 18, 23\}.$$

Each new block initially has one raw candidate. Its one hypothetical swap would create an additional raw key, so the prospective test rejects it. Both new blocks fix X , hence $F_{\text{new}}(X) = X$. The frozen old ordered rule instead gives

$$F_{\text{old}}(X) = \{-119, -113, 0, 19, 24\}.$$

The checker also confirms that the old edge and phase products separately fail to square to the identity on this same X . This explicit distinction rules out treating the old products as full-shift matchings and merely assigning them smaller radius bounds.

An independent tiny counterexample shows why isolation alone is insufficient. On the binary ring of period 12, take one type with supports

$$P = \{-1, 0\}, \quad Q = \{-1, 1\}, \quad b = r = 1, \quad H = 4.$$

Require exactness in radius one, but omit prospectivity. The input occupied sites $\{0, 1, 2\}$ move to $\{0, 1, 3\}$, which the next isolation-only update leaves fixed. The first update creates a conflicting key; the update is not an involution. The prospective rule rejects the problematic move and remains involutive. Exhaustive periodic enumeration rediscovers this mutation counterexample rather than assuming it as an expected output.

97 Executable evidence and its limits

97.1 Reference interpreter and checks

`scientific/parallel_particles.py` implements finite-support versions of both new blocks. Its separately written coordinate-local routine queries only the declared neighborhood. Before/after comparisons use complete type-anchor key sets and discard only endpoint orientation. Verification raises explicit exceptions for failures of mass conservation, candidate-set invariance, selected-set invariance, selected-orientation reversal, and block involutivity. These checks remain active under Python optimization.

The supported construction entry point accepts the frozen validated JSON source interface. Its compiler and block records are immutable under ordinary Python attribute/container operations; finite supports require exact built-in sets or frozensets of exact integers, and runtime flags require exact Booleans. Custom predicates accepted by the lower-level generic helper must independently satisfy the read-radius, purity, and endpoint-symmetry contract. The source-uniform theorem uses only the frozen compiler's validated guards. Deliberate low-level reflection or monkey-patching is outside the ordinary immutability claim.

97.2 Finite results

The primary suite compares against an independently generated micro-path and checks:

- 8928 admissible doubled states across 12 source/input cases, including 8450 states in complete reflected cycles and the entire 5670-state clean-target cycle;
- left/right increments and decrements, zero-update branches, stuck homes, disjoint guarded mergers, exact clocks, and all free/contact equality cases;
- 89 direct comparisons with the frozen old eager CA on valid states;
- all 4096 subsets of two specified eleven-site malformed universes, block verification, and whole-map inverse checks;
- 100 seeded noisy or multihead inputs, 100 translation checks, and 2074 comparisons with the declared-radius coordinate-local routine;
- a configuration with two simultaneously selected edge moves and two simultaneously selected phase moves, plus the malformed cascade above.

The independent periodic suite exhausts 24 576 binary words across three small families, with the same number of coordinate-local oracle comparisons. It checks candidate and selected-key invariance, mass, and involution. Nontrivial families include 520 words with multiple simultaneous moves, a contextual guard, and dense malformed backgrounds. The deliberately conflicting two-type family acts as identity, a limitation explicitly disclosed in its receipt. The public-interface suite checks 58 malformed-input or mutation rejections and ordinary immutability.

All suites are replayed normally and with `python -0`. The receipt files record exact case details. These finite tests are corroboration, not exhaustive enumeration of the astronomical source-uniform neighborhood, and neither the proof nor the audits are proof-assistant formalizations.

97.3 What transfers to arithmetic certificates

On a computation certified to remain in the admissible encoding, equality of old and new trajectories transports its *semantic statement*: the same encoded source run, observer event, time, or clean target is true of F_{new} . This transport must explicitly invoke the admissible-state theorem and the original certificate’s hypotheses.

Report 19’s lazy evaluator still implements the old ordered rule on arbitrary malformed finite inputs. It cannot automatically be relabeled as an evaluator of this new rule. Likewise Report 20’s ordered-factor event-budget certificate and arbitrary-malformed-state transition fibers do not become new-rule certificates by replacing a radius or a title. The malformed counterexample proves that a blanket identification would be false.

A literal new-rule arithmetic verifier must encode its own raw-key isolation, hypothetical-swap key preservation, and simultaneous update, or restrict itself to a proved admissible semantic transport theorem. This report does not provide that compiler, enumerate its fibers, or establish cheap arithmetic gate counts. Fewer composed CA blocks and a smaller sufficient spatial radius alone imply no template-independent running-time or Diophantine-witness-size bound.

[write] Report 19 is Part [IV](#) and Report 20 is Part [V](#). The two missing pieces this subsection names were supplied the same day by the next two Reports of this Part: Report 27’s sparse new-rule evaluator (Theorem [101.1](#)) and Report 28’s new-rule arithmetic compiler, which encodes raw-key isolation, hypothetical-swap rediscovery and the simultaneous update at each fixed horizon (Theorem [111.1](#)).

98 Relationship to prior work

The architecture of parallel reversible local rewrites and products of involutions is established prior art. Four primary sources are especially close.

Boyle, Lind, and Rudolph describe marker automorphisms of shifts of finite type using nonoverlapping $MDM \mapsto M\pi(D)M$ replacements in Section 2, page 74 [\[22\]](#). Their Theorem 2.6, pages 76–77, uses a product of two marker involutions. This precedes both marker-based parallel permutations and a two-involution composition as general ideas. Our candidate family has no separately installed marker word; its raw-key prospectivity is the particular local safety criterion proved here.

Toffoli and Margolus describe conserved-landscape permutations in Section 5.3 and local conservative block permutations in Section 5.5 of their review [\[23\]](#). Their recognition-preservation principle is a direct conceptual predecessor: the inverse must recognize the same local operation. The binary flip example itself is not number-conserving. We cite section numbers because the inspected author-uploaded reprint includes later corrections and annotations.

Gajardo, Kari, and Moreira characterize time-symmetric cellular automata as exactly products of two CA involutions, Proposition 1 [\[24\]](#). The algebraic time-reversal identities in Section [94](#) are an instance of that characterization, not a new architecture. Full reversibility alone does not justify an arbitrary two-block factorization; our two individual involution proofs are the substantive extra hypotheses.

Maldonado, Moreira, and Gajardo give a universal time-symmetric number-conserving CA [25]. Their Proposition 2 constructs a globally reversible conservative extension of a reversible partitioned CA by invariant recognized left, right, and malformed classes. The radius-one pair update $\overline{A}_f$ fixes its malformed class W . This statement concerns that pair update, not the complete composition $\overline{A}_f \circ \overline{I}$. Their Theorem 1 and corollaries yield the stated time-symmetric conservative universality, with an expanded alphabet; the associated time-reversing involution need not itself conserve numeric mass.

The specific result of this report is the complete-key prospective-isolation lemma with radius $3(b+r)$ and its exact application to the pinned binary five-particle compiler, including the admissible boundary and guard proof and explicit new radius. Both factors here individually conserve the binary particle number. These distinctions describe what is proved; they do not establish global priority. No novelty or optimality conclusion is drawn from this focused literature check. The primary-source audit and links are frozen with the packet.

[write] **Further lineage (added in the write; not checked at page level).** Two further traditions bear on this section. Report 26’s primary-source audit checked Kari’s two block-permutation papers at abstract level only and did not consult the other works named here. Reversible partitioned cellular automata and block representations: Morita and Harao [27]; Kari’s representation of reversible cellular automata by block permutations [28, 29]; Durand-Lose’s representation by reversible block cellular automata [30]. Number-conserving cellular automata: Boccara and Fukś [31]; Durand, Formenti and Róka [32]; Moreira [33]; and Morita and Imai [3], which Part I already cites. They are named to place the construction, not as sources of the prospective test. The M2 placement’s reading is that Lemma 90.1 is an elementary specialization of the conserved-landscape principle (a simultaneous update by nonoverlapping local involutions is an involution when each preserves its own triggering landscape); its added content is the per-candidate hypothetical-swap test, the invariance of rejected keys’ eligibility, and the explicit constants $H = 2(b+r)$ and $3(b+r)$. Report 26’s audit is shipped as `26-parallel-involutions-references-prior-art-audit.md`.

99 Portable package and verification

The companion ZIP contains this editable LaTeX source and rendered PDF; the sealed scientific packet, including independent audits, normal and optimized producer receipts, and the old compiler; inherited compiler/schema and universal-ledger references; the primary-source attribution audit; and standard-library-only release verification and replay scripts.

From an extracted release directory, run

```
python3 -B verify_release.py
python3 -B replay.py -output /an/existing/parent/new-output-directory
```

The output directory must be new and outside the release. Replay verifies the mandatory seal before running code, makes independent fresh normal and optimized copies, places logs and receipts outside those copies, and checks exact inventory, bytes, modes, and modification times after every executed script. Python bytecode writing is disabled. Scientific suites execute with the interpreter’s isolated mode and site packages disabled. No upstream repository checkout, network request, software installation, or universal source execution is needed.

The exact-inventory manifest and mandatory source pins are complementary. The inventory rejects added or missing files and directories, symlinks, unsafe paths, duplicate JSON keys, and changed bytes. The hardcoded scientific pins prevent a missing scientific source from being silently accepted after merely resealing an edited inventory. Negative tamper tests exercise

missing and changed dependencies, altered pins, malformed seals, extra files/directories, symlinks, unsafe output locations, and source-mutation detection. The PDF build command writes into a separate build directory; TeX and Poppler are optional only for rebuilding and visual checking the already supplied PDF.

The seal is an internal consistency check, not an external digital signature. An attacker able to replace both the verifier and the seal can forge a self-consistent package. The final archive digest must therefore be compared through a trusted delivery channel. Replay also does not turn bounded tests into a proof of the universal statement; that role belongs to the arguments in this report and the frozen audits.

100 Further questions

Several concrete next steps follow from the present result.

1. Can type-specific support geometry reduce $H = 2(b + r)$ or avoid paying the entire contextual guard radius for candidates that cannot compete? Any reduction must still control cooperative raw-key births and changes in initially rejected eligibility tests.
2. Can a lazy evaluator enumerate all potentially competing type-anchor keys and all prospective changes without materializing the template array? Completeness of the candidate search, especially on malformed multihead inputs, is the central obligation.
3. Can one compile the new local selection mechanism into an explicit arithmetic certificate with source-uniform size bounds? The certificate must specify whether it covers arbitrary full-shift neighborhoods or only admissible encoded computations.
4. Can the general lemma and the signed corridor partition be formalized in a proof assistant? The important distinctions are type identity versus orientation, predecessor anchor versus current head anchor, and read windows versus actual write supports.
5. Which other conservative finite-particle simulations admit the same singleton-key property on their intended encodings? The general lemma alone does not guarantee that intended dynamics survive the stronger filter.

[write] (3 October 2026.) Question 2 is answered by Report 27 (Theorem 101.1): complete sparse discovery of every competing type-anchor key, on malformed and multihead supports too, with at most $n + 2$ raw-discovery calls per step and no template array. Question 3 is answered at each fixed horizon by Report 28 (Theorem 111.1): a degree-four certificate on arbitrary finite supports of a fixed mass, malformed ones included, with source-uniform resource ledgers; its witness arity grows affinely with the horizon, so no unknown-horizon certificate is supplied. Questions 1, 4 and 5 remain open. For Question 1, a randomized probe run at placement on small synthetic families found no counterexample with H reduced to $2(b + r) - 1$ or to $b + r$; that is weak evidence, not a proof.

[write] (4 October 2026.) Question 1 is answered affirmatively for the literal templates by Part VII. Report 70 (Theorem 128.1) recognizes candidates by a guard-free geometric predicate of radius a , so that the contextual guard radius r is paid only by the local control, and isolates at $H = \max\{2(b + a), b + r\}$, which is $b + r$ for the compiler; the edge block's radius falls from $3(b + r)$ to $b + a + H$ and the sufficient radius of the whole rule from $180D + 258 + 9J$ to $108D + 149 + 3J$. Report 71 (Theorem 140.1) shortens triple exactness to b in both blocks and reaches $76D + 105 + 3J$. Both control cooperative recognition births and the eligibility of rejected anchors, as the question requires, but each does so for a new full-shift rule that differs from F_{new} on malformed inputs; neither lowers the isolation distance of Report 26's own typed-key rule. The randomized probe above is therefore consistent with these results but neither proves nor is proved

by them; whether Report 26's own rule stays an involution at $H = b + r$ is open (Section 149, item 5). Report 70 leaves further type-dependent reductions of the common exclusion bound open (Section 139).

These are open development directions here, not completed improvements. The completed conclusion is a rigorously defined binary full-shift replacement whose two individually conservative involutions preserve the entire intended five-particle dynamics with the explicit sufficient radius (95.1).

Sparse evaluation of a two involution particle cellular automaton

Research Report 27 Exact finite support evaluator and reproducibility packet

3 October 2026

Abstract of Report 27, as delivered. We give an exact finite-support evaluator for the new binary, number-conserving cellular automaton of Report 26. Its two whole-family involutions are evaluated by complete sparse discovery, all-type isolation, prospective preservation of raw keys, and simultaneous rewriting. The evaluator agrees with the eager new-rule interpreter on every finite support, including malformed configurations. Compilation stores source and explicit class metadata, without allocating the potentially enormous endpoint family. An n -particle forward or inverse step uses at most $n + 2$ raw-discovery calls, or $2n + 4$ with optional verification; there is no ordered cascade. We prove discovery completeness and give operation, bit, and space bounds, exposing the quadratic-in-cutoff class expansion. On the frozen 122,622-control source, the evaluator executes 128 literal startup steps and 128 malformed cases although the symbolic family has 269,291,358,255 types. These are bounded checks, not a completed universal startup, a Turing-machine step, or a halting result. No new Diophantine compiler or arithmetic gate count is supplied.

[write] Delivered as `Sparse_Parallel_Particle_Evaluation_Package.zip` (batch 82, archive 18; arrival db37d18c8), dated 3 October 2026; files placed by 7d2b1b245 with the prefix 27-sparse-parallel-. Report 27 evaluates the rule of Report 26 (Section 89ff.) and answers its second question. Its Lemma 103.1 restates Report 26's Lemma 90.1 with a condensed proof, and its malformed witness is Report 26's (notes at both places). The 32,034,272-byte universal source and the two universal traces it ships are not shipped here; the README rebuilds them byte for byte. Report 27's macros `\code` and `\C` are `\file` and `\calC` here; nothing printed changed. Its statements were numbered 1, 2, ... through the Report; here they are numbered by section, and its appendix is an ordinary section.

101 Result and semantic boundary

Report 26 replaces the old ordered local-factor dynamics by a *new* rule $\Phi = P_{||} \circ E_{||}$. Each displayed block is an involution on the full binary shift $\{0, 1\}^{\mathbb{Z}}$, preserves finite particle number, and commutes with the unit shift. The new and old rules agree on the entire admissible doubled source micrograph, but differ on malformed inputs. This report supplies an implementation of the new rule for arbitrary finite supports $X \subset \mathbb{Z}$, not just promised valid encodings.

The implementation is `scientific/sparse_parallel.py`. The inherited immutable source metadata, arithmetic endpoint lookup, incidence lists, and guard tables are reused; legacy ordered execution is not used. In particular, `compiler.step` evaluates the new rule, whereas `compiler.metadata.step` is an inherited *old-rule* method and is not the supported execution entry point here. The ledger field `factors` counts endpoint types, not sequential operations of the new rule.

Theorem 101.1 (Exact sparse finite-support evaluation). *For every source accepted by the frozen source schema and every finite set of integer positions X , the supported sparse step equals the eager new-rule map $\Phi(X)$. The inverse option equals $\Phi^{-1}(X)$. No endpoint array or ordered type cascade is required. If $n = |X|$, each step makes at most $2 + 2\lfloor n/2 \rfloor \leq n + 2$ raw-discovery calls; optional verification increases this to at most $4 + 4\lfloor n/2 \rfloor \leq 2n + 4$.*

The full-shift CA theorem is inherited from Report 26 and restated below. The implementation theorem concerns finite-support execution. Neither finite testing nor a roundtrip check alone proves equality of different evaluators. The proof in Sections 104–106 establishes that equality; the tests are corroboration. All earlier delivered reports remain unchanged.

102 Source data and endpoint geometry

102.1 Accepted sources and expanded metadata

The schema `reversible-two-counter-v1` represents a finite reversible two-counter partial transition system. It contains an ordered list of m distinct controls, designated start and halt controls, cutoff $J \geq 0$, and an ordered list of branches. A branch has a unique name, source and target control, side $v \in \{-1, 1\}$ identifying counter 0 or 1, update $\delta \in \{-1, 0, 1\}$, and a Boolean guard built from true, equality, greater-than, conjunction, disjunction, and negation. Every atom has a threshold $k \geq 0$ with $k \leq J$ and, on the updated counter, $k + \delta \leq J$. Enabled updates cannot make a counter negative; domains at each source and exact images at each target are pairwise disjoint. The halt has no exits.

Let p count nonzero-update branches and a count zero-update branches, in their order-preserving partitions; put $t = p + a$. The implementation validates guards on the classes $0, \dots, J, > J$, materializing domain and image tables with $(J + 2)^2$ entries per branch. Thus “source-sized” includes *expanded class descriptors*. It does not mean polynomial time or space in the binary input length of J .

The construction constants and type counts are

$$\begin{aligned} D &= 2m + 4p, & S &= 2D + 2, & B_2 &= D + 1, \\ L &= 3D + 4, & b &= B_3 = 4D + 5, & Z &= 10b + 10 + 2J, \\ N_E &= p(4D + 15) + a, & N_P &= p(4D + 14) + m, \\ F &= N_E + N_P = 8pD + 29p + m + a. \end{aligned}$$

Here L is the pair travel/exactness constant, not a common exactness radius for all templates. A template with support parameter B_g has its own exactness radius $\ell_g = 3B_g + 1$: $\ell_g = L$ for pairs and $\ell_g = 12D + 16$ for triples.

102.2 Literal endpoint families

Controls have indices $q = 0, \dots, m - 1$ and moving branches $e = 0, \dots, p - 1$. The signed mode gaps are

$$\begin{aligned} d(H_q, +) &= 2q + 1, & d(H_q, -) &= 2q + 2, \\ d(O_e, +) &= 2m + 4e + 1, & d(O_e, -) &= 2m + 4e + 2, \\ d(I_e, +) &= 2m + 4e + 3, & d(I_e, -) &= 2m + 4e + 4. \end{aligned}$$

All gaps lie in $[1, D]$, are distinct, and odd means plus. Write

$$A(M, \sigma, h) = \{h, h + d(M, \sigma)\}, \quad T(M, \sigma, h) = \{0\} \cup A(M, \sigma, h).$$

Each row below gives a pair of endpoint sets at the invariant gate anchor zero. Translating both by u gives anchor u . In moving modes $M = O_e, I_e$, set $w = v$ for outbound and $w = -v$ for inbound.

Family	First endpoint	Second endpoint
Free	$A(M, +, 0)$	$A(M, -, w)$
Behind	$T(M, +, wk)$	$T(M, -, w(k + 1)), S \leq k \leq L$
Ahead	$T(M, +, -wk)$	$T(M, -, -w(k - 1)), S + 1 \leq k \leq L$
Dispatch	$T(H_{q_s}, +, S)$	$T(O_e, -, vS)$

Family	First endpoint	Second endpoint
Endpoint	$T(O_e, +, -vS)$	$\{v\delta\} \cup A(I_e, -, v\delta - vS)$
Commit	$T(I_e, +, vS)$	$T(H_{q_e}, -, S)$
Direct	$T(H_{q_s}, +, S)$	$T(H_{q_e}, -, S)$
Phase free	$A(M, +, 0)$	$A(M, -, 0)$
Phase near	$T(M, +, \eta k)$	$T(M, -, \eta k), \eta = \pm 1, S \leq k \leq L$
Phase home	$T(H_q, +, S)$	$T(H_q, -, S)$

Free, behind, ahead, dispatch, endpoint, commit, and direct are E types. The three phase families are P types. Dispatch and direct carry the branch domain guard; commit carries its exact image guard; the other types have no class guard. Pair templates have $B_g = B_2$, all triples have $B_g = b$. The two endpoints of a type are distinct and are not translates of one another. All endpoint particles lie in $[-b, b]$.

Arithmetic lookup in `gate_at(i)` decodes the existing ordered partition into these families without scanning preceding types or a travel interval. Moving E types come in blocks of $4D + 15$: two travel modes of $2D + 6$ types each, then dispatch, endpoint, commit. Direct types follow. Moving P modes each have $2D + 7$ types: phase free and $2(D + 3)$ near phases; home phases follow. E identities are $0, \dots, N_E - 1$; P identities are offset by N_E . These integers are stable *names*, not an execution schedule.

For reference, a valid home encoding is

$$\{-Z - c_0, 0, Z + c_1, S, S + d(H_q, \sigma)\}, \quad c_0, c_1 \in \mathbb{N}.$$

The admissible set also includes enabled-edge intermediate encodings and all translations. The sparse algorithm never assumes its input lies in this set.

103 The new parallel rule

Put $I_s(u) = [u - s, u + s] \cap \mathbb{Z}$ for $s \geq 0$, and let the block read bound satisfy $r \geq b$. The type-specific write interval is $W_g = I_{B_g}(0)$, translated to $u + W_g$. A type-anchor key (g, u) is raw when the particles in $I_{\ell_g}(u)$ are exactly one translated endpoint of type g , and its contextual guard holds. The orientation $\lambda \in \{0, 1\}$ is the endpoint label, stored separately. Distinct types at one anchor are distinct keys.

A guard reads the two intervals

$$[u - Z - J, u - Z], \quad [u + Z, u + Z + J].$$

An empty interval represents class $J + 1$, a singleton represents its exact signed offset from distance Z , and two or more particles reject the guard. Its table is the branch domain table or, for commit, the exact image table. These sites lie outside the gate write interval and are unchanged by an endpoint swap.

For a block, write $\mathcal{C}(X)$ for its complete raw-key set. Let τ_k be the local involution exchanging the two equal-weight endpoint words on the type's write interval and fixing other words. For a raw key $k = (g, u)$, its hypothetical support $Y_k = \tau_k X$ is obtained by deleting the current endpoint particles and adding the other endpoint particles. Set $H = 2(b + r)$. Select k precisely when

1. $k \in \mathcal{C}(X)$;
2. no other key of $\mathcal{C}(X)$ has anchor within distance H of u ;

3. $\mathcal{C}(Y_k)$ and $\mathcal{C}(X)$ have identical keys anchored within distance $b + r$ of u .

Condition 3 ignores orientation but retains *every* raw type, including keys failing isolation or prospectivity. All decisions use the same X ; only then are all selected rewrites performed simultaneously. E and P use separate within-family tests, with

$$r_E = Z + J, \quad r_P = 3b + 1, \quad H_E = 2(b + r_E), \quad H_P = 2(b + r_P).$$

The old raw radius B_g , old same-type exclusion, and ordered refresh-and-continue evaluator are not used.

Lemma 103.1 (Prospective isolation, Report 26). *For finitely many equal-weight local endpoint swaps of write radius at most b with translation-covariant, endpoint-symmetric candidate predicates of read radius at most r , the above selected simultaneous map is a full-shift, translation-equivariant, particle-conserving involution of radius at most $3(b + r)$.*

Proof. Selected anchors are more than H apart, so their writes are disjoint. A swap at u can affect a raw predicate at v only for $|v - u| \leq b + r$. Thus prospectivity is equality of the entire raw-key set after that one swap. A predicate read window cannot meet two selected write windows, since that would put their anchors at distance at most $2(b + r) = H$. Every predicate therefore has its original truth value after simultaneous writing, ruling out cooperative candidate births.

The prospective condition of a key reads within radius $b + 2r$ of its anchor. Other selected writes are outside this window by separation. Its own swap interchanges the two sets being compared and is involutive; consequently every isolated key, including one rejected by prospectivity, retains its prospective status. The preserved raw-key set preserves isolation, hence the selected-key set is preserved. Reapplying the same disjoint swaps restores the input. Isolation reads radius $H + r$ around a key, while writing a fixed coordinate requires an anchor within b . The output radius is $b + H + r = 3(b + r)$. Shift covariance and equal-weight disjoint swaps give the remaining conclusions. $\square$

[write] **A restatement.** Lemma 103.1 is Report 26's Lemma 90.1, and its proof is Report 26's proof in condensed form; it is kept as Report 27's own presentation, not as a second route, and no novelty is claimed for it. The radii displayed next are Report 26's (Section 95, equation (95.1)).

The type-specific write interval is important: unrelated particles outside a pair's own write interval must not be erased merely because they lie inside the common bound b . Applying the lemma to E and P gives

$$\begin{aligned} R_E &= 3(b + Z + J) = 33b + 30 + 9J, \\ R_P &= 3(b + 3b + 1) = 12b + 3, \\ R &\leq R_E + R_P = 180D + 258 + 9J. \end{aligned}$$

Report 26 also proves agreement with the old dynamics on every admissible doubled micrograph state, including travel-boundary equalities, branch image guards, and reflections. That preservation theorem is bundled unchanged with its two audits. It is not a claim of full-shift equality to the old rule.

104 Complete sparse candidate discovery

104.1 Recovering every possible type

Sort the finite support and enumerate every occupied close pair $(h, h + d)$ with $1 \leq d \leq D$. The gap uniquely identifies mode and sign. Every endpoint contains such a pair; a triple contains

exactly one. Indeed its marker is separated from the head by at least S , and from the second pair particle by at least $S - D = D + 2 > D$. This remains true at the reverse endpoint interaction, where marker and pair have both shifted. The whole input may have many heads or close pairs: none are discarded.

For every moving close pair, always emit the free and phase-free types. For every distinct third particle z , write $\rho = h - z$ and $\epsilon = 0$ for plus, 1 for minus. Emit the types specified by the following exhaustive tests:

$$\begin{array}{ll} \text{behind:} & k = w\rho - \epsilon, & S \leq k \leq L, \\ \text{ahead:} & k = -w\rho + \epsilon, & S + 1 \leq k \leq L, \\ \text{phase near:} & k = |\rho|, \quad \eta = \text{sign}(\rho), & S \leq k \leq L. \end{array}$$

The four moving interaction cases are

$$\begin{array}{ll} O_e, + & \rho = -vS \quad \text{endpoint,} \\ O_e, - & \rho = vS \quad \text{dispatch,} \\ I_e, + & \rho = vS \quad \text{commit,} \\ I_e, - & \rho = -vS \quad \text{endpoint.} \end{array}$$

For a home close pair H_q , if $h - S \in X$, emit its home phase and every outgoing dispatch/direct incidence when plus, or every incoming commit/direct incidence when minus. Source incidence lists are stored once, with $O(t)$ total entries.

Lemma 104.1 (Containment completeness). *If either translated endpoint of type g is contained in X , the sparse type-discovery procedure emits g , without any admissibility, exactness, guard, or isolation assumption.*

Proof. The endpoint's close pair is enumerated and determines its signed mode. A pair endpoint is covered by unconditional free/phase-free emission. For a triple, its distinct marker is enumerated. Substitution into the endpoint table yields exactly the displayed behind, ahead, phase, interaction, or home test, including all equality boundaries. A behind minus endpoint has $\rho = w(k + 1)$, and an ahead minus endpoint has $\rho = -w(k - 1)$, explaining the sign corrections. A reverse endpoint interaction retains $\rho = -vS$ even though its marker has moved. Home incidences include every matching branch at that control. These cases exhaust both orientations of every family. $\square$

The procedure only proposes identities. It does not assume that the current left pair coordinate is the invariant anchor. This distinction is essential both for a free reverse endpoint and for the moved marker of a reverse interaction.

104.2 Recovering every raw key and orientation

Filter proposed identities to the requested family. For each type, materialize its two sorted shapes Q_0, Q_1 . For each orientation and every occupied site z , try

$$u = z - \min Q_\lambda.$$

Retain $(g, u) \mapsto \lambda$ exactly when all $u + Q_\lambda$ are occupied, the closed interval $I_{\ell_g}(u)$ contains exactly $|Q_\lambda|$ particles, and the guard holds. Two binary searches count an interval in sorted X . The same interval operation implements guard classes without scanning the $J + 1$ detector sites.

Proposition 104.2 (Raw-map equality). *The sparse map contains exactly the eager new-rule keys with the same orientations on every finite X .*

Proof. A genuine raw occurrence has a contained endpoint, so Lemma 104.1 supplies its type. Its least occupied endpoint site is tried, giving exactly its invariant anchor. The remaining literal endpoint, inclusive exactness, and guard tests are the definition of raw occurrence. Conversely every returned key passes those same tests. Distinct endpoint sets cannot both equal the exact occupied interval at one key, so orientation is unambiguous. $\square$

Optional center/radius arguments filter *returned anchors* inclusively. They do not crop the support used for discovery or guards. This prevents losing candidates whose defining particles lie outside the requested anchor interval. Prospective rediscovery is performed from scratch on the hypothetical support, rather than reusing the original proposed identities.

105 Selection equality and local execution

Sort all raw keys by (u, g) . A key has another key within distance H if and only if its predecessor or successor in this order is within H . Thus the nearest-neighbor test is exactly the all-type isolation test, even when different types share the same anchor.

For each isolated key, swap its endpoints on the unchanged original X , discover the complete new raw map, and compare keys anchored within $b + r$ before and after. Keys are compared without orientation, and the implementation additionally checks that the retained key's orientation reversed. Rejected candidates remain in the original raw map as competitors. Once all decisions are fixed, delete/add the selected endpoints simultaneously.

Proof of Theorem 101.1, equality part. Proposition 104.2 gives identical original and hypothetical raw maps. Hence nearest-neighbor isolation, prospective key equality, and selected key/orientation maps agree with the eager new-rule interpreter. Lemma 103.1 makes selected writes disjoint, so the final support agrees exactly. Applying block equality first to E then to P proves forward equality. Applying it to P then E proves inverse equality because both blocks are involutions. No argument assumes a valid encoding or five particles. $\square$

Each block also exposes a coordinate-local output oracle. To determine output at i , it uses only $X \cap I_{3(b+r)}(i)$. Writing anchors lie within b of i ; competitors lie within H of those anchors, and raw predicates read a further r . Thus every required coordinate lies within $b + H + r = 3(b + r)$. Prospective reads also fit within this crop. Its candidate queries remain complete for the relevant anchors. This API is a literal local rule oracle; its repeated raw calls are not covered by the whole-step call count below.

A useful malformed witness is $X = \{-118, -112, 0, 18, 23\}$ for a two-control right increment source with $J = 0$. Each block initially has one raw key, but its hypothetical swap creates another raw key. Prospectivity rejects both moves, so the new rule fixes X . The inherited old rule maps it to $\{-119, -113, 0, 19, 24\}$. This distinguishes the implemented maps rather than merely showing reversible behavior.

[write] This witness is Report 26's (Section 96); Report 28's certificate accepts (X, X) and rejects the old rule's image on it (Section 120.3).

106 Work space and bit bounds

106.1 A bound before prospective rejection

Let $n = |X|$. All keys admitted to the hypothetical-swap stage are mutually more than $H \geq 2b$ apart. Their original endpoints occupy disjoint neighborhoods and each uses at least two particles of the same X . Their number is therefore at most $\lfloor n/2 \rfloor$. This counts *all isolated keys*, not just the ultimately selected ones, so rejected probes are included.

Exactness ensures each hypothetical swap replaces an equal-sized endpoint without colliding with outside particles; it preserves n . The final simultaneous block also preserves n , so both block inputs and all discovery snapshots in one step have mass n . Each unverified block performs one initial raw discovery plus one for every isolated key. This proves the $2 + 2\lfloor n/2 \rfloor$ bound. Tracing does not repeat selection. With verification, each block additionally discovers the final support and selects once more, then checks the reverse write using that computed selected map. This doubles the call bound to $4 + 4\lfloor n/2 \rfloor$. Empty and singleton supports are included.

106.2 Expanded source and runtime costs

Let Δ bound the longest home-in or home-out list; $\Delta \leq t$. Let G be the total guard-program/snapshot size, $c = (J + 2)^2$, and M denote immutable metadata size including payloads. Compilation stores

$$O(m + t + G + tc) \text{ words, plus actual string and integer payloads.}$$

Class validation takes expected $O(cG + tc + m + t)$ unit operations, with parsing/name costs charged as well. Dense growing conflict masks have a deliberately loose additional $O(tc^2)$ bit-work bound. Large binary-encoded J can therefore be expensive, even though subsequent sparse guards do not scan detector intervals site by site.

For one step let C bound the number of discovered identities over all original, hypothetical, and verification supports, including both families before filtering. Close pairs, third particles, and incidence emissions give

$$C = O(n^3 + n^2\Delta), \quad Q = 1 + n^3 + n^2\Delta.$$

The sharper cap $C \leq F$ is also valid. If K bounds raw keys in one family, trying at most $2n$ anchors per type gives $K \leq 2nC$. Under expected constant-time hash access and unit-cost exact integer arithmetic, one raw call takes

$$T_{\text{raw}} = O(n \log(n + 1) + Q + C \log(C + 1) + Cn \log(n + 1)).$$

Actual name-string copying is charged separately: each materialized template copies only its source names and the needed index digits. The implementation sorts original raw keys, scans at most K keys per isolated probe for its before-set, and constructs one hypothetical support at a time. A safe expected step bound is

$$O((n + 2)T_{\text{raw}} + K \log(K + 1) + nK + n^2),$$

with a constant-factor increase for verification. Substituting $C = O(Q)$ and $K = O(nQ)$ gives a polynomial in particle count and explicit source/class parameters without an unconditional F loop or coordinate-span scan. Dense inputs can discover all F types; no worst-case sublinear-in- F claim is made.

Work storage beyond M is safely bounded by $O(n + Q + nC)$ words. The original raw map, one hypothetical map and restricted key sets, and the output support coexist; previous hypothetical

maps are discarded. A step trace contains at most n selected triples. There is no persistent materialized-template cache, which could otherwise grow to the entire type family over many steps.

106.3 Integer lengths and deterministic bounds

Let $A = \max_{x \in X} |x|$, or zero for empty X . Since equal-sized endpoints lie in $[-b, b]$, an arbitrary particle matching moves at most $2b$ per block, hence at most $4b$ over two actual blocks. Hypothetical and verification snapshots obey the same order of magnitude. Anchors and detector reads are bounded by $A + O(b + Z + J)$, with no ordered- F displacement term.

Let W include the bit lengths of that coordinate bound, F , J , source numeric values, and container indices; let L_{name} bound source-name lengths. Addition, comparison, and integer hashing cost $O(W)$ in a simple bit model; the bounded integer multiplication/division in decoding can be charged $O(W^2)$. The expected-hash claim is not a worst-case constant-time claim about Python dictionaries.

For an intentionally loose deterministic bound, put $U = 1 + M + n + Q + nC$, dominate a container operation by at most U key comparisons, and charge $O((W + L_{\text{name}} + 1)^2)$ per scalar operation. Multiplying the displayed runtime operation bound by $U(W + L_{\text{name}} + 1)^2$ supplies a conservative polynomial bit bound in the expanded parameters. It is not tight or a practical performance estimate, and it does not remove the non-polynomial-in-log J class expansion. Compilation's separate mask and parsing costs still apply.

107 Supported API and reproducibility contract

Create `a = SparseParallelCompiler(source)` from a validated decoded JSON object. Use `a.step(X)` or `a.step(X, inverse=True)`. `verify=True` checks mass, candidate-key invariance, selected-key invariance, reversed orientations, and each block's involution. `trace=True` returns the final support and immutable step/block records. `a.E` and `a.P` expose candidates, eligibility, application, and local output. `a.gate_at(i)` materializes one template by stable global ID, and `a.encode(q, c0, c1, sign)` creates home encodings.

Position inputs must be exact built-in sets or frozensets whose elements are exact Python integers; Booleans, floats, and integer subclasses are rejected. Flags must be exact Booleans. Empty supports and arbitrary positive or negative integer coordinates are valid. Validation uses explicit exceptions, so it survives `python -O`. The source snapshot, compiled records, returned discovery/selection maps, and traces are immutable through ordinary attribute/container operations. Mutating the caller's source after compilation cannot change semantics. Low-level reflection and monkey-patching are outside this immutability contract. The fresh ordinary dictionary returned by `ledger()` may be modified without affecting the compiler.

Only source-derived immutable class guards are supported. Arbitrary caller-supplied executable predicates are not a public construction path. The bundled inherited source document describes its older API; the enclosing `SparseParallelCompiler` and the entry points above define this report's API.

108 Measured evidence and independent checks

108.1 Small-source differential testing

The frozen main suite executes 8,183 configurations over 10 sources and checks 1,140 arithmetic/eager template equalities. It records 16,366 comparisons/checks in each of the following categories: raw candidate maps, selected maps, block outputs, block involutions, whole-step equalities (forward and inverse combined), and whole-step roundtrips. Candidate equality includes exact invariant anchors and labels, not just final support or reversibility.

Coverage includes 2,280 endpoint orientations, 2,280 endpoint-plus-spoiler cases, 3,072 exhaustive malformed subsets over 11-site and 10-site universes, 192 encoded states, 120 random supports, 20 huge translations, 1,761 local-radius comparisons, and 162 guard-detector configurations. There are 259 API/mutation rejections and 19 source-validation parity checks. Thirty-two instrumented probe-bound cases include empty and singleton supports, prospective rejection, competition, valid states, and three separated copies. Maximum observed discovery counts are 8 without verification and 16 with it, within the mass-dependent bounds.

A 780,035-type source is compiled while eager construction and old-execution entry points are trapped. Compilation constructs zero endpoint templates; bounded runtime queries touch only five distinct types. The frozen main-suite elapsed times are 54.252 seconds normally and 50.512 seconds optimized. These machine-specific observations are not performance guarantees.

An independently authored all-index small-source oracle additionally checks 14 varied source cases, 17,504 contained-endpoint discoveries, 1,120 full raw key/orientation maps, 2,240 block candidate/selection/output comparisons, and 2,560 sparse/literal guard comparisons. Its final producer replay times are 3.796 and 4.156 seconds. The independent symbolic/code audit found no correctness defect; it is not a proof-assistant formalization.

108.2 Frozen universal-source benchmark

[write] The source that Report 27 ships as `scientific/universal-source.json` is Report 17's (SHA-256 `fa61d06178d7...`, Section 76). It and the two universal traces of this benchmark are not shipped in this directory; the README gives the commands that rebuild all three byte for byte from shipped files in a few seconds.

The included 32,034,272-byte source has

$$(m, p, a, J) = (122622, 66066, 75495, 0), \quad D = 509508,$$

$$F = 269291358255, \quad R = 91711698.$$

The benchmark disables old ordered execution and eager construction before metadata compilation. It runs 128 *literal new-rule* startup steps with inverse verification and exact comparison against the frozen old admissible-state trace. It then runs 128 arbitrary endpoint/noise configurations sampled across the type range, including translation by -2^{2048} , with mass and roundtrip verification. No source-boundary jump is substituted for a CA step.

Frozen producer measurement	Normal	Optimized
Metadata compilation, seconds	3.914	7.033
128 literal startup steps, seconds	0.136	0.194
128 malformed cases, seconds	0.118	0.168
Total benchmark, seconds	4.532	7.853
Peak process RSS, KiB on Linux	376,776	378,012

The bounded runs observe at most four candidate types, two raw keys, two selected keys, and two prospective probes per block. Optimization’s longer compilation in this particular run reflects ordinary run variation, not different semantics. The benchmark imports the standard-library `resource` module, and its reported RSS units are Linux KiB.

These checks do not finish the long universal startup, execute a whole simulated Turing-machine step, establish halting, or materialize an eager universal oracle. The malformed benchmark supplies roundtrip corroboration; arbitrary-support equality rests on the proof and independent small-source comparisons. The frozen source’s attributed universal construction is inherited rather than newly proved by this report.

108.3 Frozen packet and release verification

The published scientific subtree contains 39 files: 37 manifest entries plus its manifest and checksums. All scientific code, data, proofs, tests, and benchmark receipts are unchanged. Four inert metadata files are adapted for portable provenance: `PROVENANCE.json` uses relative packet origins; `audit-timings.json` names its interpreter as `python3`, retaining every flag and measured value; `manifest.json` and `SHA256SUMS` are regenerated. The other 35 files preserve their original bytes, modes, and modification times. Universal traces remain byte-identical to the producer.

The unchanged independent root review included with the references records proof/code inspection, both interpreter-mode replays, semantic receipt comparisons, and invariance of all 39 files in the *original* scientific packet. It does not claim to review the later publication metadata or this article. The original scientific manifest is preserved as a pinned reference; its hash is distinguished from the published adapted manifest below. No executable scientific semantics changed.

The completed release adds exact-inventory verification and mandatory hardcoded source pins. Omitting the universal fixture is an error, not a reduced replay mode. Negative tests operate on disposable copies and exercise missing or added entries, unsafe metadata and paths, altered pins and dependencies, and symlinks. A byte-level privacy scan rejects known internal absolute-path prefixes; portable temporary-output examples remain allowed. Release replay uses six fresh relocated copies, one for each suite/mode combination, with external outputs, bytecode disabled, isolated Python, standard-library dependency checks, and a network audit guard. It checks the exact file and directory inventory, bytes, modes, and modification times after every subprocess.

Receipt comparison normalizes only declared timing/RSS fields and validated known optimization flags. All trace/hash/count fields are retained; universal trace bytes are compared exactly against the producer. A seal proves internal consistency, not authenticity against replacement of the verifier and all pins. The final delivery hashes provide an external reference.

109 Attribution and arithmetic certification limits

The two-block architecture belongs to established marker-automorphism and conserved-landscape traditions. Boyle, Lind, and Rudolph describe marker-delimited permutations and marker involutions [22]; Toffoli and Margolus discuss conserved landscapes [23]. Recognition-set preservation in Maldonado, Moreira, and Gajardo, Proposition 2, is a close proof precedent [25]. Their construction uses a different state space and does not itself supply this binary two-conservative-involution compiler. The equivalence between CA time symmetry and a product of two CA involutions is established in Gajardo, Kari, and Moreira [24]; reversibility alone is not that assertion.

The bundled bounded primary-source audit records exact checked passages and access limits. It does not establish priority, radius minimality, or global novelty. This report’s contribution

relative to the frozen construction is complete source-metadata-based discovery and exact sparse execution, together with the explicit probe and resource bounds.

An old arithmetic verifier for the ordered rule does not automatically certify the new map on malformed supports. Semantic equality transports admissible trajectories, but arbitrary new-rule certificates must encode complete raw discovery, all-type exclusion, hypothetical key-set preservation, and simultaneous writes. Two mathematical involutions do not mean two cheap arithmetic instructions. This release contains no new Diophantine compiler, no new universal equation, and no revised arithmetic gate count.

Useful next questions are to compile these exact finite-support predicates into a separately audited arithmetic verifier; account for sorting, integer encodings, and complete candidate witnesses; reduce explicit guard-class expansion where sound; and measure longer literal trajectories with clearly stated stopping conditions. They remain future work.

[write] (3 October 2026.) The first of these questions is answered at each fixed horizon by Report 28 (Theorem 111.1), whose emitter has an independent code audit and charges sorting (a bitonic network), signed integer encodings and complete candidate witnesses. Guard-class reduction and longer literal trajectories remain open.

110 Offline replay and artifact provenance

[write] This section is Report 27’s Appendix A. The full replay requires Linux, Python 3.10 or newer, the standard library, and permission to create symlinks for negative tests. The exact tested interpreter is recorded in the replay receipt. No network, external sibling packet, or package installation is required. Allow roughly 370 MiB peak memory for the frozen benchmark, with headroom for process variation. Run from the extracted release directory:

```
python3 -B verify_release.py
python3 -B release_checks.py
python3 -B replay.py --output /tmp/report27-fresh-replay
```

The chosen output directory must not exist and must be outside the release. Full replay automatically runs normal and optimized variants of the main differential suite, independent candidate audit, and universal benchmark. It writes the final replay receipt outside the sealed inventory. The original frozen scientific scripts can also be run separately with their `-output-dir` option; see the release README.

Rebuilding the article is a separate authoring action using a local TeX installation and `sh build.sh`; it is not part of scientific replay. It replaces only the requested PDF and invalidates the existing release seal until the changed deliverable is reviewed and resealed. PDF rendering/inspection and exact final file hashes are recorded in the release QA and delivery provenance.

The following SHA-256 values distinguish the original scientific basis from the adapted published inventory. The release manifest pins all remaining files individually.

Original scientific manifest

311e6293d0b28528b63e8f9e0eadd1a4fa91d7dc5c1b208c3e2654be37cdd7fa

Published scientific manifest

5672555af673f129c52aed9317c4d0c68b3c31fbd926952aae385e4ab32c8528

Sparse evaluator

785082d121622f8fc755ed7039ee6efe719964296eec7d76ad44495dc5c6c2c1

Universal source

fa61d06178d718511232a80a05e02192d940fde3c273f0c202f0623253b2ade3

Report 26 scientific manifest

141f1ecf4e34f9046b738fa6caf2d3a80c4e74f46aad7d9fa4e3b140d634f26b

All scientific code, data, proofs, tests, and benchmark receipts are byte-identical local copies; only the four disclosed provenance/inventory metadata files are adapted. No upstream repository code was downloaded or executed in this release. The frozen local compiler and evaluators were executed for the stated checks. The report source, PDF, all scientific fixtures, proofs, audits, receipts, and replay tools are distributed together; no scientific dependency is optional.

A fixed horizon quartic compiler for the parallel particle rule

Report 28

3 October 2026

Abstract of Report 28, as delivered. For a fixed validated finite reversible source, a fixed particle mass n , a fixed whole-step horizon T , and a fixed time direction, we construct an ordinary integer polynomial of degree at most four that certifies the exact endpoint relation of the new parallel two-involution particle rule. The $4n$ external natural variables encode two strictly increasing lists of signed integer coordinates. Each accepted external fiber has exactly one complete natural witness; every other fiber has none. The compiler directly enumerates source records and occupied-particle slots, uses deterministic sorting and compacted hypothetical lanes, and performs full prospective candidate rediscovery. It neither enumerates the entire endpoint-template array nor bounds the coordinates. A two-particle example is delivered as a fully expanded quartic with 12,595 nonzero monomials. A five-particle malformed example distinguishes this relation from the old ordered rule. Exact syntax and coefficient ledgers, independent audits, bounded regression evidence, a complete-witness linear height bound, and a portable replay package accompany the theorem. The result is a fixed-horizon family, not a fixed-arity unbounded-time representation or a real-witness uniqueness theorem.

[write] Delivered as `Canonical_Parallel_Quartic_Certificates_Package.zip` (batch 82, archive 03; arrival `db37d18c8`), dated 3 October 2026; files placed by `7d2b1b245` with the prefix `28-parallel-quartic-`. Report 28 certifies the rule of Report 26 and uses Report 27’s evaluator only as a test oracle. It shares its compiler pattern with Report 20 (note in Section 114); its sorting network and arithmetic layer are credited to CDC and SMC Part IV in notes in Sections 115 and 117. The archived receipts `research/supplement-receipt*.json` predate the final compiler (footnote in Section 123). Report 28’s macro `\N` (printing $\mathbb{N}_0$) is `\Nzero` here; nothing printed changed. The delivered table of contents is omitted and its appendix is an ordinary section.

111 The precise result

The new cellular automaton is the composition $\mathcal{F} = P \circ E$ of two prospective-isolation blocks. These blocks use the endpoint geometries inherited from the finite reversible-source construction, but their action on arbitrary particle supports is different from the old ordered-factor scheduler. Reports 26 and 27 established and evaluated that new rule; the present report supplies its concrete fixed-horizon arithmetic compiler. The scientific dependencies needed to identify the construction are included in the release, rather than left as unavailable external references.

[write] Reports 26 and 27 are the first two Reports of this Part (Sections 89ff. and 101ff.) [20, 21]; Report 28 has no bibliography entries for them.

Let $\mathbb{N}_0$ include zero. Fix source data s accepted by the bundled source parser, $n, T \in \mathbb{N}_0$, and $\epsilon \in \{+1, -1\}$. These four parameters are compiler syntax. For natural external variables

$$\mathbf{e} = (x_0^+, x_0^-, \dots, x_{n-1}^+, x_{n-1}^-, y_0^+, y_0^-, \dots, y_{n-1}^+, y_{n-1}^-),$$

write $x_i = x_i^+ - x_i^-$ and $y_i = y_i^+ - y_i^-$. A signed pair is canonical when its product is zero.

Theorem 111.1 (Concrete exact-endpoint quartic). *The delivered compiler constructs a finite, explicitly indexed list $R_1, \dots, R_M \in \mathbb{Z}[\mathbf{e}, \mathbf{z}]$ of degree at most two, with W witness variables $\mathbf{z}$, and hence the ordinary polynomial*

$$P_{s,n,T,\epsilon}(\mathbf{e}, \mathbf{z}) = \sum_{j=1}^M R_j(\mathbf{e}, \mathbf{z})^2 \quad \text{of degree at most four.}$$

For every $\mathbf{e} \in \mathbb{N}_0^{4n}$, there is exactly one $\mathbf{z} \in \mathbb{N}_0^W$ satisfying $P_{s,n,T,\epsilon} = 0$ if and only if all external signed pairs are canonical, both coordinate lists are strictly increasing, and

$$\mathcal{F}^{\epsilon T}(\{x_0, \dots, x_{n-1}\}) = \{y_0, \dots, y_{n-1}\}.$$

Otherwise there is no such witness. Uniqueness includes every inactive, rejected, sorting, routing, and padding register.

No coordinate bound or promise of a valid source encoding is assumed. In particular, the statement covers malformed, noisy, and multihead supports of the fixed mass. It is an exact requested-endpoint relation; an observer or halting predicate is not silently supplied. The coefficient ring is $\mathbb{Z}$, but the quantification of both external variables and witnesses is over $\mathbb{N}_0$. No uniqueness assertion is made over $\mathbb{R}$, $\mathbb{Q}$, or arbitrary signed auxiliary integers.

The theorem is not a single polynomial of fixed arity for an unknown horizon. Witness arity grows affinely with the fixed T ; source and mass also affect syntax. It does not imply a finite-fold representation of arbitrary recursively enumerable sets or settle a finite-fold MRDP question. Earlier observations about infinite native Pell witness fibers are therefore not contradicted: this is a different, finite-horizon, deterministic circuit substrate.

[write] The “earlier observations about infinite native Pell witness fibers” are uncited in Report 28. They are Report 22’s: its canonical clean-clock certificates for the three-mass automaton have infinite native extension fibers. Report 22 is source 23 of **signal-machine-collision-certificates** (Part VIII, batch 82M2). (Added 3 October 2026, batch 82, reciprocal note: its fixed-scale infinitude is Theorem `smc:ch:thm:infinite` there, Section 155.)

112 Finite source and endpoint geometry

112.1 Validated source data

A source has an ordered list of m controls, designated start and halt controls, a nonnegative class cut J , and an ordered finite branch list. A branch has a source control, target control, side $v \in \{-1, +1\}$, update $\delta \in \{-1, 0, +1\}$, and a finite Boolean guard over two natural counters. Guard atoms test equality or strict comparison with nonnegative integer constants; conjunction, disjunction, negation, and truth are supported. The left side updates counter 0 and the right side counter 1.

The parser checks the literal JSON schema, exact field types, nonnegative post-counters, pairwise-disjoint branch domains at each source, and pairwise-disjoint exact images at each target. The halt control has no outgoing branch. Domain and exact-image guards must be constant on the two-counter classes $0, \dots, J, > J$; the parser explicitly builds their $(J + 2)^2$ Boolean tables. The image table is derived, including pre-counter nonnegativity, rather than accepted as unverified supplied data. This is a finite partial-injection source contract. Its complete literal specification is `research/references/SOURCE_SCHEMA.md`.

Let p be the number of nonzero-update moving branches and a the number of zero-update direct branches. Set

$$\begin{aligned} D &= 2m + 4p, & S &= 2D + 2, & L &= 3D + 4, \\ B_2 &= D + 1, & b = B_3 &= 4D + 5, & Z &= 10b + 10 + 2J. \end{aligned}$$

For control index i , moving-branch index e , and sign $\ell \in \{0, 1\}$ (plus, minus), the distinct head gaps are

$$g_H(i, \ell) = 2i + 1 + \ell, \quad g_O(e, \ell) = 2m + 4e + 1 + \ell, \quad g_I(e, \ell) = 2m + 4e + 3 + \ell.$$

Thus all signed gaps are distinct in $1, \dots, D$, plus gaps are odd, and minus gaps are even. Write $[h; g] = \{h, h + g\}$ for a head pair. All displayed endpoint coordinates below are relative to their invariant anchor.

112.2 The complete endpoint families

For a moving branch e , let its source and target controls be i, j . In a moving mode $M \in \{O, I\}$ put $w = v$ for O and $w = -v$ for I , and abbreviate $g_M^\ell = g_M(e, \ell)$. The E endpoints are as follows; the left and right columns are paired orientations 0 and 1.

Family	Orientation 0	Orientation 1
Free	$[0; g_M^0]$	$[w; g_M^1]$
Behind	$\{0\} \cup [wt; g_M^0]$ $S \leq t \leq L$	$\{0\} \cup [w(t+1); g_M^1]$
Ahead	$\{0\} \cup [-wt; g_M^0]$ $S+1 \leq t \leq L$	$\{0\} \cup [-w(t-1); g_M^1]$
Dispatch	$\{0\} \cup [S; g_H(i, 0)]$	$\{0\} \cup [vS; g_O(e, 1)]$
Endpoint	$\{0\} \cup [-vS; g_O(e, 0)]$	$\{v\delta\} \cup [-vS + v\delta; g_I(e, 1)]$
Commit	$\{0\} \cup [vS; g_I(e, 0)]$	$\{0\} \cup [S; g_H(j, 1)]$
Direct	$\{0\} \cup [S; g_H(i, 0)]$	$\{0\} \cup [S; g_H(j, 1)]$

Dispatch, endpoint, and commit occur once per moving branch; direct occurs once per direct branch. Free, behind, and ahead occur for both moving modes. Dispatch uses the branch domain guard, commit its exact-image guard, and direct its domain guard, in *both* endpoint orientations. All remaining families are unguarded.

The P families just change a signed head gap, without moving its head anchor or marker:

Family	Orientation 0	Orientation 1
Phase free	$[0; g_M^0]$	$[0; g_M^1]$
Phase near	$\{0\} \cup [st; g_M^0]$	$\{0\} \cup [st; g_M^1]$
Phase home	$\{0\} \cup [S; g_H(i, 0)]$	$\{0\} \cup [S; g_H(i, 1)]$

Phase near occurs for each moving mode, side $s \in \{-1, +1\}$, and $S \leq t \leq L$; phase home occurs for each control. Pair families have write radius B_2 ; all triple families have radius b . Their displayed endpoints are contained in those intervals. In a triple the marker is farther than D from both head particles: the lower separation bound is $S - D = D + 2$. The endpoint interaction moves marker and head together, so preserves these relative separations. Consequently every endpoint contains exactly one close pair of gap at most D .

For a type g of write radius B_g , set $\lambda_g = 3B_g + 1$. At anchor u its raw predicate requires that all particles in the *closed* interval $[u - \lambda_g, u + \lambda_g]$ are exactly one of its two translated endpoints. If guarded, inspect the closed bands

$$[u - Z - J, u - Z], \quad [u + Z, u + Z + J].$$

An empty band has class $J + 1$; a singleton has its distance from the inner endpoint as class; multiple particles make the guard false. The fixed table is then applied to the two classes. These contextual bands lie outside the write interval. Raw matching is therefore endpoint-symmetric. Pair exactness uses $3B_2 + 1 = L$, not the triple radius and not the old radius- B_g raw test.

113 The new parallel rule and its inverse

Write $\mathcal{C}_B(X)$ for the set of raw keys (g, u) in block $B \in \{E, P\}$. Orientation is a property of an occurrence, not part of the key. Distinct types at the same anchor are distinct competitors. Set

$$r_E = Z + J, \quad r_P = 3b + 1, \quad q_B = b + r_B, \quad H_B = 2q_B.$$

The raw predicate at u reads within radius r_B . A key $k = (g, u)$ is selected exactly when it is raw, no other raw key has anchor distance at most H_B , and its hypothetical single rewrite preserves the entire raw key set whose anchors are within q_B of u . All selected endpoint rewrites are made simultaneously on the immutable input X . Raw-but-prospectively-rejected keys still participate in isolation.

Lemma 113.1 (Prospective isolation). *Each such block is an everywhere-defined, translation-equivariant involution on binary configurations. On finite supports it conserves particle number. Its radius is at most $3(b + r_B)$.*

Proof. Selected anchors have distance greater than $2(b + r_B)$, so their radius- b write sets are disjoint. A rewrite at u can alter a raw predicate at v only if $|v - u| \leq b + r_B$. Hence the prospective test is equivalent to preserving the global raw key set under that single rewrite. A raw-predicate read window meets at most one selected write set: if it met two, their anchors would be at distance at most $2(b + r_B)$. It follows that simultaneous rewrites preserve every raw predicate, excluding cooperative births as well as individual births.

Raw-set invariance preserves isolation. The prospective test for a key at u depends only on radius $b + 2r_B$ about u . Every other selected write is outside that region, by the strict anchor separation. The key's own rewrite preserves its prospective status because it is an involution and interchanges the two raw sets compared. Thus all isolated keys, including rejected ones, retain their status. The selected set is unchanged, and a second application reverses the same disjoint swaps. Equal endpoint cardinalities give finite-particle conservation. Isolation reads radius $H_B + r_B$ about a key; an output site can be written only from within b . Therefore an output radius is

$$b + \max\{H_B + r_B, b + 2r_B, r_B\} = 3(b + r_B).$$

All constructions commute with translations. Disjoint finite writes also define the block on infinite configurations, although the certificate below addresses only finite supports. $\square$

Thus $\mathcal{F} = P \circ E$ has inverse $E \circ P$. A sufficient whole-rule radius is

$$R_{\text{new}} = 3(b + Z + J) + 3(4b + 1) = 180D + 258 + 9J.$$

This is inherited rule geometry, not a claim about small arithmetic circuit cost. The all-input block proof and the valid-encoding preservation proof are also included as pinned dependencies. In particular the old and new rules agree on the admissible doubled source micrograph, but need not agree elsewhere. Section 120.3 makes the distinction explicit.

114 Direct source slots and exact type identifiers

The arithmetic compiler does not materialize the full endpoint-template array. Fix increasing occupied coordinates $x_0 < \dots < x_{n-1}$. Enumerate occupied pairs $h = x_i < h_2 = x_j$ for $i < j$, a distinct occupied marker $z = x_k$ when needed, source record, moving mode, and orientation ℓ . A travel distance is computed from coordinate differences and range-tested; it is never enumerated through $S, \dots, L$.

[write] **The same compiler pattern as Report 20.** Reports 20 and 28 share the source schema and endpoint geometry, the index constants (Report 28's e_s, e_m, p_m, ρ below are Report 20's A_E, U_E, V_P, r of Section 78), the occupied-pair and marker slot enumeration (Report 20's (80.1); here (114.1)–(114.2)), the guard bands, the simultaneous displacement formula (Report 20's (81.1)), the arithmetic gadgets and the mass-two example source. They differ in the rule and in the theorem: Report 20 certifies the old ordered rule under an event budget ($T + K$ scheduler rounds, $2n$ external inputs, a mass-two exporter only), Report 28 the new parallel rule at a fixed horizon in either direction ($4n$ external inputs, a general emitter measured at masses 0–5). Neither proves the other's theorem, and Section 120.3 shows why an ordered-rule certificate cannot be relabelled.

For exact type IDs put

$$e_s = 4D + 15, \quad e_m = 2D + 6, \quad p_m = 2D + 7, \quad \rho = D + 3,$$

$$E_{\#} = pe_s + a, \quad P_{\#} = 2pp_m + m.$$

For moving branch e and mode number $k = 0$ for O , $k = 1$ for I , put $e_b = ee_s + ke_m$ and $p_b = E_{\#} + (2e + k)p_m$. IDs are global across E then P .

Family	Type ID	Anchor and recovery formula
Free E	e_b	$u = h - \ell w$
Behind	$e_b + 1 + t - S$	$u = z, t = w(h - z) - \ell, S \leq t \leq L$
Ahead	$e_b + 1 + \rho + t - (S + 1)$	$u = z, t = -w(h - z) + \ell, S + 1 \leq t \leq L$
Dispatch	$ee_s + 2e_m$	$u = z; h - z = S$ at $\ell = 0, vS$ at $\ell = 1$
Endpoint	$ee_s + 2e_m + 1$	$u = z - \ell v \delta; h - z = -vS$ in both orientations
Commit	$ee_s + 2e_m + 2$	$u = z; h - z = vS$ at $\ell = 0, S$ at $\ell = 1$
Direct	$pe_s + d$	$u = z, h - z = S$; direct branch index d
Phase free	p_b	$u = h$
Phase near	$p_b + 1 + \sigma \rho + t - S$	$u = z, t = s(h - z), S \leq t \leq L; \sigma = 0, 1$ for $s = -1, +1$
Phase home	$E_{\#} + 2pp_m + i$	$u = z, h - z = S$; control index i

Every row additionally tests its appropriate signed gap from Section 112. It constructs the old and new endpoint lists, counts all n coordinates in the type-specific exactness interval, and requires that count to be the endpoint arity. Because the old coordinates were chosen from the occupied list, presence is already known. A false travel-range test substitutes a fixed in-range dummy before forming the ID while keeping the enabling predicate false. Rejected endpoint coordinate expressions are still fully defined; no dynamic unchecked array lookup is used.

A guarded slot scans all n particles in each detector band, computes hit count and singleton offset, and then scans all $(J + 2)^2$ table cells with equality and Boolean gates. False cells are also charged. For a non-singleton the class is deterministically $J + 1$; a count above one also forces rejection. No invalid class is used as an array index. Domain versus image is fixed by type, independently of orientation.

The exact widths are

$$K_E = \binom{n}{2} [4p + (14p + 2a) \max(n - 2, 0)], \quad (114.1)$$

$$K_P = \binom{n}{2} [4p + (8p + 2m) \max(n - 2, 0)]. \quad (114.2)$$

Per pair, two modes and two orientations give $4p$ pair slots. Per distinct marker, E has $8p$ behind/ahead slots, $6p$ interaction slots, and $2a$ direct slots. P has $8p$ near slots and $2m$ home slots. The implementation checks these widths explicitly.

Lemma 114.1 (All-input raw binding). *On every strictly increasing integer support, the live direct slots are exactly the raw key map of the indicated block, with no duplicated live key.*

Proof. Each true endpoint has the unique close pair proved in Section 112. In a triple the remaining marker is therefore unique. Its signed gap identifies mode and orientation; the formulas above recover travel, side, and invariant anchor. Every true endpoint is generated by the corresponding one source/pair/marker/orientation slot. In particular reverse free motion uses $u = h - w$, and reverse endpoint interaction uses $u = z - v\delta$; replacing either by the current head or marker would give the wrong key.

Conversely, a live slot passes the literal gap, incidence, travel, exactness, and guard conditions, so is a genuine raw occurrence. For a fixed type, anchor, and orientation the unique close pair and marker exclude duplicate generation. The two distinct equal-size endpoint sets cannot both be the entire support in the same exactness interval. Hence opposite orientations cannot duplicate the same key. Each family is emitted only once. This argument uses neither a unique global head nor any valid-encoding assumption. $\square$

115 Sorting isolation and deterministic lane compaction

For $K > 0$, pad the direct slot list to the next power of two N . Sort records (live, u , slot) by a fixed bitonic network with live records first and increasing signed anchor within either live status. The padding record is $(0, 0, -1)$. Every comparator and every selected payload field is computed, including those in inactive rows. Exact comparison ties do not swap at that comparator; global stable sorting is unnecessary because the entire network wiring is fixed.

The sorting statements here use standard comparison-network correctness, not the finite sorting test corpus. A monotone threshold map commutes with compare-exchange, so a network that sorts all zero-one inputs sorts every totally ordered input. In the bitonic construction, sorting two half-blocks in opposite directions makes a bitonic sequence; the distance-half compare layer partitions it into two bitonic halves with every lower-half value at most every upper-half value. Recursive merging proves correctness by induction on the power-of-two width. The alternating adjacent-comparison network with n phases is the classical odd-even transposition sorting network; its all-input sorting theorem, together with the zero-one principle and bitonic network theorem, is treated in Knuth, *The Art of Computer Programming*, volume 3, second edition, Section 5.3.4 [26]. Payload fields follow the same deterministic swaps, so neither theorem requires arbitrary tie permutations.

[write] The deterministic comparator with payload and the explicit bitonic network, with its correctness proof and unique comparator auxiliaries, are printed in `canonical-diophantine-certificates`, Proposition `cdc:mem:prop:comparator-cost` and Subsection `cdc:mem:subsec:bitonic`; Report 28 cites Knuth [26] only. No novelty is claimed for them.

Lemma 115.1 (Neighbor isolation). *For a live key in the sorted live prefix, another raw key is within distance H if and only if an immediately adjacent live record is within distance H .*

Proof. The live prefix is in nondecreasing anchor order. Any nonadjacent neighbor on either side is at least as far away as the adjacent one on that side. Equal-anchor distinct types are adjacent to some equal-anchor key and conflict at distance zero. Inactive rows cannot separate live rows. Lemma 114.1 ensures a key does not falsely compete with a duplicate of itself. $\square$

Isolation is computed against the complete raw prefix before prospectivity. Prefix sums of isolated bits in this sorted order assign deterministic ranks. There are exactly $\lfloor n/2 \rfloor$ fixed lanes whenever slots exist. Lane r scans every sorted record to obtain the source-slot index with isolated rank r , then scans every direct slot to route its descriptor. Missing lanes have presence zero, index -1 , and all-zero ID, anchor, arity, old sites, and new sites. A pair's third padded site is zero, but its arity is two, so this padding cannot act as a real particle at coordinate zero.

Lemma 115.2 (Lane capacity). *The fixed lane count never discards an isolated key on a genuine support.*

Proof. Isolated anchors have separation greater than $H = 2(b+r) \geq 2b$. Their endpoint supports lie in disjoint radius- b intervals, and each uses at least two of the n particles. There can be at most $\lfloor n/2 \rfloor$ such keys. $\square$

No arbitrary permutation, free rank, free absent-lane descriptor, or existential sorting choice is introduced. Each is a function of previously computed wires. On rejected external inputs these same formulas still give a total deterministic circuit, even where geometric capacity or distinctness need not hold.

116 Prospective rediscovery and simultaneous transport

For each lane, replace its old endpoint by its new endpoint when present, and use the original support otherwise. Transport is implemented by checking each original particle against each old site and adding the corresponding displacement. Sort the result with n fixed odd-even transposition phases, using strict greater-than swaps. Live single swaps have distinct coordinates; the comparator network is nevertheless fully defined on ties as well.

Every lane, including an absent lane, runs a complete direct raw circuit on its hypothetical support. No initially missing type is excluded from this pass. For an isolated lane key $k = (g, u)$, the complete before-key set in the closed $q = b + r$ neighborhood is just $\{k\}$, since $q \leq H$. Thus full prospective equality is equivalent to

1. some after-key has the same type ID and invariant anchor as k ;
2. no different after-key has anchor within q of u .

The comparison deliberately ignores orientation. Endpoint symmetry ensures the own key survives with the opposite orientation, while the explicit own-key test binds this fact in the emitted circuit. Checking only initially discovered types would miss the obstruction in Section 120.3.

The final block transport again starts from the same immutable input. For each original particle x , its new coordinate is

$$x + \sum_{\text{lanes } \lambda} \sum_{j=0}^2 b_{\lambda} [j < \text{arity}_{\lambda}] [x = o_{\lambda j}] (n_{\lambda j} - o_{\lambda j}),$$

with products lifted through quadratic gates. Here b_λ is the selected bit, and o, n are old and new sites. Selected old supports are disjoint. Exactness prevents collision with unchanged particles, and anchor separation prevents collisions between selected endpoints. The fixed head/head/marker pairing need not preserve individual common sites: its image set is still exactly the endpoint replacement. A final deterministic coordinate sort gives the canonical output.

Proposition 116.1 (Whole-step circuit binding). *On every canonical strictly increasing support, an emitted block computes the corresponding new parallel block. Composing E, P for T iterations computes $\mathcal{F}^T$; composing P, E computes $\mathcal{F}^{-T}$.*

Proof. Lemma 114.1 binds discovery. Neighbor isolation and lane capacity bind all isolated keys. The singleton-before argument and full after rediscovery bind the prospective test. Simultaneous transport binds the selected set replacements. Mass and distinctness are preserved, so canonical sorting supplies the same hypotheses to the next block. Induct over the $2T$ blocks, using Lemma 113.1 for the inverse order. $\square$

117 Quadratic arithmetic and complete witness uniqueness

Every wire is either a natural bit or a signed value represented by natural coordinates. Constants are literal integers in polynomial expressions, not existentially selectable values. A nonlinear intermediate is assigned to a fresh register before nonlinear reuse; all defining residuals have degree at most two. The implementation rejects a higher residual degree with an explicit exception.

117.1 Signed assignment

For a fully specified integer expression f in earlier wires, introduce $p, m \in \mathbb{N}_0$ and residuals

$$p - m - f = 0, \quad pm = 0.$$

The only possible values are $(p, m) = (\max(f, 0), \max(-f, 0))$. If $f = 0$, both vanish. The conclusion holds for every integer value of f , without a magnitude bound. Here f has degree at most two, so both residuals are quadratic or lower.

117.2 Exact integer comparison

For signed integers x, y already determined, put $d = x - y$. Introduce $b, s \in \mathbb{N}_0$ and impose

$$b(b - 1) = 0, \quad d - (2b - 1)s + 1 - b = 0. \quad (117.1)$$

The first equation gives $b = 0$ or 1 . If $b = 1$, the second gives $s = d \geq 0$. If $b = 0$, it gives $s = -d - 1 \geq 0$, which is possible exactly when $d < 0$. The cases are disjoint and exhaustive over the integers, including $d = 0$ and $d = -1$. Hence $b = [x \geq y]$ and its slack are uniquely forced. Nonnegative slack, not a nonvanishing inverse gadget, enforces the strict alternative.

Equality is the bit

$$[x = y] = \text{ge}(x, y) - \text{ge}(x, y + 1).$$

Strict order and closed intervals use these comparisons. Every comparator pays two fresh natural variables and two residuals.

117.3 Boolean operations and selections

For previously proved bits, xy , $x + y - xy$, and $1 - x$ implement AND, OR, and NOT. A derived Boolean register $z = f$ uses the one residual $z - f$; no extra Booleanity residual is needed because induction proves $f \in \{0, 1\}$. Empty AND and OR folds are the constants one and zero. A conditional integer selection is the signed assignment

$$x + b(y - x),$$

which returns x for $b = 0$ and y for $b = 1$. It is a definition for both cases, not an implication guarded by b . Multiplication by a false mask never deletes the defining equations of the masked calculation.

Lemma 117.1 (Total topological uniqueness). *For each natural external assignment, before imposing pure acceptance checks there is exactly one complete natural assignment to all auxiliary registers.*

Proof. Traverse the fixed operation order. Signed assignment has exactly one canonical natural pair for every already determined integer expression. Comparison has exactly one natural bit/slack pair by (117.1). Boolean gates receive already determined bits and have exactly one bit output. Constants, folds, selects, prefix sums, sorting payloads, failed raw descriptors, absent lanes, and full rejected hypothetical calculations all use these same total definitions. No branch removes a defining equation. Induction therefore determines every auxiliary coordinate and provides an existing natural value for each. The argument applies even when an external input will subsequently fail canonicity or ordering checks. $\square$

[write] **Credit and a second presentation.** Signed assignment, the comparison gadget (117.1), equality as a difference of comparison bits, Boolean gates and total topological uniqueness are the canonical circuit layer of `canonical-diophantine-certificates` (Lemmas `cdc:pt:lem:signed`, `cdc:pt:lem:circuit` and `cdc:mem:lem:compare`) and of SMC Part IV (`smc:sl:eq:cmp`, `smc:sl:eq:eq`); Report 20 proves the same layer (Lemma 79.1). Report 28 cites none of them. This section is kept as Report 28's own presentation; no novelty is claimed for it.

117.4 Acceptance checks and the quartic

The pure checks are $2n$ external pair products, $2(n - 1)$ strict-order acceptance rows when $n \geq 1$, and n final coordinate equalities. Order comparison auxiliary variables are included in the topological circuit. A check may appear early in the emitted row list; logically it filters the unique gate assignment just proved.

Proof of Theorem 111.1. By Lemma 117.1, any zero of all defining residuals has the same complete natural auxiliary assignment. External pair products and order checks enforce exactly the canonical domain. On that domain, Proposition 116.1 identifies the computed endpoint with the requested iterate. The endpoint rows retain the unique assignment exactly when the target list equals that endpoint. Finally a finite sum of integer squares vanishes exactly when every summand vanishes. Squaring quadratic residuals gives degree at most four and does not change the natural zero set. This proves existence, soundness, completeness, and complete-witness uniqueness. $\square$

This is an algebraic proof, not an assumption that the supplied witness evaluator is correct. A supplied assignment can be checked directly against the residuals or expanded quartic without

invoking that evaluator. One-coordinate mutation tests provide useful binding evidence but cannot replace Lemma 117.1.

118 Resources that are actually charged

Let A count signed assignment registers, I comparison bit/slack pairs, G derived Boolean output registers, and Q pure acceptance rows. The exact identities are

$$W = 2A + 2I + G, \quad M = W + Q, \quad Q = \begin{cases} 5n - 2 & n \geq 1, \\ 0 & n = 0. \end{cases} \quad (118.1)$$

There are $4 \max(n - 1, 0)$ order-comparison witnesses. If w_E, w_P are the emitted witness counts for one block, excluding domain and endpoint work, then in either direction

$$W(T) = 4 \max(n - 1, 0) + T(w_E + w_P). \quad (118.2)$$

This is an exact affine horizon-arity formula, not a fixed-arity claim.

The counter A measures signed-register assignments, *not* uniformly bounded-fan-in elementary additions. Count, offset, rank, and displacement expressions can be wide affine sums. They must be charged by residual syntax, monomial occurrences, and coefficient sizes. A circuit ledger therefore records all of

$$\sum_j |\text{supp } R_j|, \quad \sum_{j,\mu} \text{bits}(|c_{j,\mu}|), \quad \sum_{j,\mu} |c_{j,\mu}|, \quad \max_{j,\mu} \text{bits}(|c_{j,\mu}|),$$

where only nonzero coefficients count and $\text{bits}(k) = \lfloor \log_2 k \rfloor + 1$ for $k > 0$. These magnitude-bit sums exclude sign and index serialization overhead; exact file bytes are reported separately where a file was emitted.

For $c = (J + 2)^2$ and a nonzero block slot width K , one raw pass has conservative cost $O(n^2 + K(n + c + 1))$ in construction work and sparse residual monomial occurrences. The bitonic isolation network costs $O(K \log^2(K + 1))$; routing costs $O(nK)$. There are $\lfloor n/2 \rfloor$ full hypothetical raw passes, each preceded by $O(n^2)$ transport/sorting work. Final transport and sorting are $O(n^2)$. Thus a conservative complete-block bound is

$$O(K \log^2(K + 1) + nK(n + c + 1) + n^3). \quad (118.3)$$

A zero-slot block is semantically identity, but for $n \geq 2$ the present emitter still assigns one signed pair-gap register per occupied index pair. Those $O(n^2)$ assignments remain in its ledger; they must not be optimized away in reporting. For $n < 2$ even this scratch is absent. Sum the appropriate E and P bounds across T whole steps and add $O(n)$ external checks. This concrete sorting bound supersedes the earlier design audit's quadratic-isolation macro estimate.

Bound (118.3) concerns residual syntax, not the monomial count after expanding the ordinary quartic. The exact number of ordered product-term occurrences before combining like terms is

$$B_{\text{SOS}} = \sum_j |\text{supp } R_j|^2 \quad (118.4)$$

This is an upper bound on the number of distinct expanded monomials. Cancellation can change the final monomial and coefficient totals. The widest residual has $O(n + 1)$ terms, giving a further conservative factor $O(n + 1)$ when converting the residual-syntax bound to an expansion bound. The concrete pair example is fully expanded; larger examples were emitted and tested as explicit finite residual systems in memory, not claimed to have been stored as expanded polynomials.

No coefficient hides a packed endpoint-template array. The number of geometric types is $E_{\#} + P_{\#} = 8pD + 29p + m + a$, but production compilation loops over source records and occupied slots instead of those travel-indexed types. Source validation still pays explicit finite guard tables; there is no polynomial-in-log J claim and no claim of practical universal-source construction.

There is no coordinate bound in the polynomial syntax. If a genuine initial support has $\max_i |x_i| = A_0$, particle transport is at most $2b$ per block, so genuine coordinates through T whole steps have magnitude at most $A_0 + 4bT$. Slacks, rejected descriptors, and other scratch values are still charged witness variables and may have input-dependent bit length. Coefficients of the fixed-source, fixed- n, T polynomial do not grow when its external coordinate assignment is translated.

119 An explicit bound on every witness coordinate

The following companion result is about the actual emitted operation families. Quadratic residual degree alone would not suffice: a generic quadratic assignment circuit could repeatedly square a growing value. Here every product involving coordinate-dependent data is a Boolean mask times an affine expression; the other products are between bounded controls. Constraint products and the final SOS squares do not allocate new value registers.

[write] In this section H_* is an isolation distance, $\max(H_E, H_P)$. Report 20's witness-height bound, which Report 20 also calls H_* , is printed Ht_* (Section 85).

Let $\ell = \lfloor n/2 \rfloor$, $F = E_{\#} + P_{\#}$, and let P_E, P_P be the least powers of two at least $\max(1, K_E)$ and $\max(1, K_P)$, respectively. These P_B are *padded occurrence widths*, not endpoint-type counts. Define the explicit, deliberately conservative constants

$$\begin{aligned} H_* &= 2[b + \max(Z + J, 3b + 1)], \\ \alpha &= 14 + 12\ell, \\ \beta &= (10 + 6\ell)b + Z + J + 2H_* + n(J + 1) + F + P_E + P_P + 10, \\ C(s, n) &= (4b + 1)(\alpha + \beta). \end{aligned}$$

Proposition 119.1 (Complete natural witness height). *Fix s, n . If an external endpoint pair is accepted and $A_0 = \max_i |x_i|$ (zero at empty mass), then every one of its unique auxiliary natural coordinates z_j obeys*

$$0 \leq z_j \leq C(s, n)(A_0 + T + 1).$$

With zero represented by one bit, the sum of ordinary binary component lengths is at most

$$W(T) \lceil \log_2(C(s, n)(A_0 + T + 1) + 2) \rceil = O_{s, n}((T + 1) \log(A_0 + T + 2)).$$

This includes inactive lanes, rejected raw/prospective scratch, padding, and both external sortedness scans.

Proof. For one block whose genuine input coordinates have magnitude at most R , a raw descriptor has old-site magnitude at most R , new-site magnitude at most $R + b$, and anchor magnitude at most $R + 1$, even when its raw bit is false. Travel masking bounds IDs by F ; counts are at most n , ranks and indices are bounded by the padded widths, and arities by three. Detector flags restrict each nonzero offset term to $[0, J]$, so offset sums are at most nJ . Direct inspection of the gap, incidence, exactness, detector and table comparisons, including equality's extra unit shift, bounds every raw witness on magnitude- U inputs by

$$\text{RawCap}(U) = 2U + 3b + Z + J + n(J + 1) + F + 10.$$

Signed natural components are no larger than the represented absolute value; comparison slacks are no larger than their integer difference magnitude. Boolean outputs are at most one. Sorting and routing select existing values, so repeated selects do not amplify their semantic magnitudes.

A single hypothetical lane has three padded displacement terms, each at most $2R + b$ in magnitude. Hence its transported coordinates are at most $7R + 3b$, and its sorting slacks at most $14R + 6b + 1$. Full raw rediscovery is bounded by $\text{RawCap}(7R + 3b)$. The prospective comparisons have slack at most $8R + 3b + H_* + 3$ or $F + 1$. Final transport has at most 3ℓ displacement terms; its coordinates and sorting slacks are bounded by

$$(1 + 6\ell)R + 3\ell b, \quad (2 + 12\ell)R + 6\ell b + 1.$$

Initial sorting/routing comparison slacks are bounded by $2R + H_* + 3$ or $P_E + P_P + n + 2$. Every displayed linear coefficient and constant is dominated by α and β , respectively. This covers a complete block, including every disabled operation, by $\alpha R + \beta$ without assuming disjointness in hypothetical scratch.

For the accepted computation the actual geometric block boundaries and accepted target satisfy $R \leq R_* = A_0 + 4bT$. Apply the one-block bound separately at each genuine boundary; do not recursively amplify a coarse hypothetical-scratch bound across steps. The external order slacks are at most $2R_* + 1$ and satisfy the same majorant. Thus

$$\alpha R_* + \beta \leq (\alpha + \beta)(R_* + 1) \leq C(s, n)(A_0 + T + 1).$$

Multiplying the component binary-length bound by the exact arity gives the claim. $\square$

The accepted-target restriction matters: an arbitrarily large wrong external target can have arbitrarily large pre-check order slack while A_0, T stay fixed. It has no accepted witness and is outside this height claim. The binary bound counts component lengths, not JSON bytes or self-delimiting encoding overhead. The complete structural audit and an independent derivation are included under `bitbound/`; finite checks inspect 66 emitted blocks, 332 complete block assignments, and 72 accepted endpoints in both interpreter modes. The source-uniform conclusion is the structural proof, not extrapolation from these examples.

120 Concrete emitted examples

120.1 A fully expanded two-particle quartic

The source in `research/example-source.json` has controls q, h , one true-guard right increment $q \rightarrow h$, and $J = 0$. Thus $m = 2, p = 1, a = 0, D = 8, S = 18, L = 28, B_2 = 9, b = 37, Z = 380$, and 95 endpoint types. The new radius bound is 1,698; the inherited ordered-source radius field is 21,144 and refers to the old rule.

At $n = 2, T = 1$ in the forward direction, the accepted endpoint is

$$X = \{0, 5\}, \quad E(X) = \{1, 7\}, \quad Y = P(E(X)) = \{1, 6\}.$$

The external natural prefix is $(0, 0, 5, 0, 1, 0, 6, 0)$. There are four direct slots and one lane per block. The exact measured ledger is

Part	A	I	G	W	M
One E block	110	167	191	745	745
One P block	110	167	191	745	745
Complete certificate	220	336	382	1,494	1,502

The complete certificate has eight pure checks and 1,502 total polynomial variables including the eight external coordinates. The coincident equality of total variables and residuals here is accidental; the general identities are (118.1).

Its 1,502 quadratic residuals have 5,511 monomial occurrences, 6,544 coefficient-magnitude bits, coefficient ℓ_1 mass 20,032, maximum coefficient magnitude length 10 bits, and ordered SOS expansion bound 29,249. The expanded ordinary quartic has

Measured quantity	Exact value
Degree	4
Nonzero monomials after combination	12,595
Coefficient-magnitude bits	29,413
Maximum coefficient magnitude bits	23
Coefficient ℓ_1 mass	6,567,618
Serialized quartic JSON bytes	220,969
Serialized residual-system JSON bytes	86,561

The residual JSON, quartic JSON, complete accepted assignment, and independent parser/expansion verifier are included. The hardened verifier checks exact schemas and integer arities, natural values, canonical monomial indices, and the binding of input/target metadata to the external natural prefix. It imports neither `compiler.py` nor `circuit.py`. It reconstructs the literal square-sum, compares the ordinary coefficient map, evaluates every residual and the quartic, and rejects all 1,494 upward single-witness-coordinate mutations. Its separate validation suite rejects 37 malformed cases in both interpreter modes; an independent audit adds 31 negative probes in each mode. It checks the supplied algebraic artifacts and metadata; arbitrary ledger contents and variable-name semantics are not independently reconstructed. CA/source semantics depend on the audited compiler and source pins. The exact expanded polynomial SHA-256 is

```
ddd6973f1285d60c2ae663b6b47cd933
827e9e86c332cb69148c1ef45c27dca1
```

120.2 Measured mass dependence at one whole step

For this same source, the actual complete forward certificates have the following sizes. All variables means $4n + W$.

n	K_E	K_P	W	M	All variables
0	0	0	0	0	0
1	0	0	0	3	4
2	4	4	1,494	1,502	1,502
3	54	48	47,080	47,093	47,092
5	460	400	784,928	784,951	784,948

At mass five, $A = 169,152$, $I = 142,492$, $G = 161,640$, and $Q = 23$. The residual system has 3,312,070 monomial occurrences, 3,826,624 coefficient-magnitude bits, coefficient ℓ_1 mass 13,669,134, maximum coefficient magnitude length 10 bits, and ordered SOS expansion bound 20,580,990. These are current implementation costs, not lower bounds, optimality claims, or numerical improvements over earlier certificates with different semantics and scope.

120.3 The malformed old versus new separator

For the same source let

$$X = \{-118, -112, 0, 18, 23\}.$$

The E raw map has the single key $(0, -119)$ in orientation 1. Its hypothetical replacement gives $\{-119, -114, 0, 18, 23\}$. Inside its prospective radius $q_E = 417$, the after map contains the own key reversed and the new key $(1, 0)$. Hence the lane is rejected.

The P raw map on X has the single key $(47, -118)$ in orientation 1. Its hypothetical replacement gives $\{-118, -113, 0, 18, 23\}$. Inside $q_P = 149$, the after map contains the own key reversed and the new key $(59, 0)$. That lane is also rejected. Thus both blocks fix X , so $\mathcal{F}(X) = X$ and $\mathcal{F}^{-1}(X) = X$.

The old ordered rule instead gives

$$Y_{\text{old}} = \{-119, -113, 0, 19, 24\}.$$

The new mass-five one-step certificate accepts (X, X) and rejects (X, Y_{old}) in both directions. This tests genuine isolated raw occurrences and prospective births, rather than a vacuous input with no candidates. It also demonstrates why an old ordered-factor arithmetic certificate cannot simply be relabeled as a certificate of the new rule.

121 Boundary cases and interface semantics

For $n = 0$ or 1 no pair slot exists; every block is identity. At $n = 0$ the polynomial is zero with no external or auxiliary variables, and the unique complete witness is the empty tuple. At $T = 0$ the compiler retains canonical signed-pair, sortedness, and identity-endpoint checks. At $p = 0$ there are no moving-mode slots; direct E and home-phase P families remain, and an empty E block is allowed.

The supported emitter constructor is `Certificate(source,n,T,inverse)`. It validates exact natural n, T and an exact Boolean direction, then calls the pinned source parser for validation and immutable metadata. It does not call a frozen transition, candidate enumerator, gate lookup, or ordered-factor evaluator. Its production circuit is new code.

The standard `witness(x,y)` helper constructs an assignment and checks every residual. By contrast, `witness(...,check=False)` and `result(x)` are unchecked diagnostics; their return is not acceptance or input validation. Definitive acceptance means a natural full assignment with every residual zero. `serialize()` produces explicit finite SOS residual syntax; `polynomial()` expands and combines ordinary monomials and can need much more memory.

122 What the finite computations establish

The all-input claims rest on the proofs above and source-level audits. Regression tests bind that mathematics to the emitted code on specified finite inputs. Tests are active both normally and under Python optimization; essential checks use explicit exceptions, not removable assertions.

122.1 Baseline and nonvacuous exhaustive domains

The baseline reports 1,066 complete raw-map equalities on 533 fixture snapshots, 694 accepted endpoint equalities, 1,388 intermediate block/selection equalities, 678 wrong-target rejections, two old/new separators, and 67,867 upward or downward natural witness-coordinate mutation

checks. It also covers canonical-pair, duplicate, unsorted, empty, singleton, and zero-horizon cases.

The exhaustive portion is precisely all mass- n subsets of each of the following four finite universes, in both whole-step directions:

Source	Coordinate universe	n	Supports
No branches	$\{0, 6, 7, 8, 11\}$	3	10
One direct branch	$\{0, 10, 11, 12, 13, 14\}$	3	20
Right increment	$\{0, 1, \dots, 8\}$	2	36
Right increment	$\{-18, -13, -10, 0, 18, 19, 23, 24\}$	3	56

There are $2(10 + 20 + 36 + 56) = 244$ exhaustive endpoint checks. Each domain contains at least one raw/selected event. In the no-branch source E is identity but home-phase P still moves two of the ten supports; it is not an entirely inert test domain. The remaining corpus is targeted, not exhaustive over all supports or source schemas.

122.2 Supplements and independent audit

The supplements add 104 endpoint and 208 block equalities, 275 all-zero absent-lane payload checks, cases with two simultaneously selected lanes, 28 accepted and 28 rejected horizon fixtures, 20 large-translation tests (including shifts $\pm 10^{80}$), 16 exact horizon-arity checks, and 56 endpoint checks with multiple source records. The direct source guards include left/right increments and decrements, distinct domain/image predicates, and guarded direct branches.

The independent code audit additionally checks 822 raw snapshots against a literal all-template test oracle, 686 live endpoint shapes, 425 sorting cases, 270 compaction/padding cases, and 1,350 full lane payloads. Of the compaction cases, 104 have multiple occupied isolated lanes. Its verdict is that no correctness blocker was found; this is a mathematical/source-code audit, not a proof-assistant certification. Literal template enumeration is confined to the test oracle, not production compiler syntax.

The horizon supplement uses all mass-at-most-three horizon fixtures through $T = 3$, together with mass-five fixtures only for $T \leq 1$. Independently generated mass-five $T = 2, 3$ CA answers are present in the fixture corpus, but *their certificate tests were not completed and are not claimed as passed*. An earlier oversized supplemental run was stopped; the bounded suites were then rerun in both modes. The mathematical horizon theorem is proved by composition, rather than inferred from uncompleted runs.

122.3 Timing and verification limits

Original baseline runs took approximately 118 seconds normally and 104 seconds optimized; the bounded supplements took approximately 92 and 81 seconds, and multisource tests approximately 9 seconds in either mode. These are environment-dependent observations, not complexity benchmarks or required deadlines. Fresh release replay timings are recorded separately in machine-readable external receipts. The release replay runs suites sequentially to avoid overstating practical capacity through simultaneous large allocations.

No universal-source-size polynomial was emitted. No finite test covers the full infinite coordinate domain, and mutation checks do not prove global uniqueness. The expanded pair verifier checks that concrete delivered polynomial; the general source-uniform theorem additionally relies on the compiler binding and topological proof. The delivered code is research software, not a formally verified implementation.

123 Release contents and independent replay

The release ID is `parallel-quartic-release-20261003`; the standalone article is Report 28. `report28.tex` rebuilds `report28.pdf` with pdfLaTeX. The scientific packet is under `research/`; it contains all production code, frozen parser/evaluator dependencies, source schema, geometry and preservation proofs, independent design and emitter audits, full mandatory finite fixtures, test generators, test runners, typed receipts, exact ledgers, and the concrete source/SOS/quartic/witness example. Nothing scientifically necessary is an optional private fixture.

`RELEASE-MANIFEST.json` specifies exact relative-path inventory, byte lengths, modes, and SHA-256 digests. `RELEASE-SHA256SUMS` provides the digest list including the manifest. `verify_release.py` rejects missing or extra files, symlinks and nonregular entries, digest or mode mismatches, source-pin mismatches, and malformed receipt schemas. It verifies the included scientific manifest and provenance in addition to the outer release inventory. The manifest's own final digest is supplied with the release receipt; a local integrity checker alone is not an external trust root.

The original scientific freeze is retained as provenance with manifest digest

```
dc2106553312d36cbcd655e5e3551e52
771682de0a1f7f27e91798561ecf90a0
```

The release incorporates a later, explicitly recorded output-routing revision to the test runners. The four test-runner changes are output-routing only. A separately disclosed repair hardens the independent example verifier schema and metadata binding, with its own negative test suite. Compiler, circuit, mathematical fixtures, and the acceptance relation retain their pinned identities.¹ Exact original and final manifest digests, modified-file scope, and historical fixture-path rebinding are recorded in provenance. Stable relative paths and report/release IDs replace ephemeral runtime locations. Reports 26 and 27 are unchanged.

123.1 Requirements and read-only procedure

Replay needs a Linux/POSIX filesystem and Python 3.10 or later with standard-library modules only; no package installation or network access is required. The mathematical emitter itself is ordinary Python; Linux is the verified release replay platform, particularly for POSIX modes and timestamp checks. Rebuilding the PDF additionally needs a TeX distribution with the packages named in the source. The PDF is already provided, so TeX is not required to verify the mathematics or software outputs.

From the extracted release directory, choose a fresh writable output directory *outside* the release and run

```
python3 -B verify_release.py
python3 -B replay_release.py --output-dir ../report28-replay
```

The replay helper uses absolute resolved paths internally, disables bytecode, runs the independent concrete verifier and all four bounded suites, the example-validation suite, the independent negative probes, and the full witness-height audit normally and with `-O`, and directs every

¹[write] One exception: the archived `research/supplement-receipt.json` and its optimized copy lack four ledger fields (`endpoint_type_count`, `legacy_ordered_radius`, `new_rule_radius`, `source_metadata_semantics`) that the pinned `compiler.py` emits, so an earlier compiler revision produced them; their counts agree otherwise. They are not shipped. The fresh replay receipt `expected/normal/supplement-receipt.json`, shipped as `data/28-parallel-quartic-expected-supplement-receipt.json`, is the binding evidence.

generated receipt/log outside the release. It compares parsed JSON with strict type checks and exact required field sets. Only explicitly designated runtime fields (elapsed seconds and the mode flag) may differ. Scientific counts, full ledgers, source data, fixture pins, and audit results must agree with the frozen references and between modes.

Before and after replay it snapshots the complete file inventory, bytes, POSIX modes, and nanosecond modification times; all must be identical. Reads may alter filesystem access times, which are deliberately not included in this claim. Source permissions are never rewritten to simulate immutability. A fresh output receipt records the actual elapsed times and release digest. A final archive replay re-extracts the ZIP before checking; delivery claims should be read together with its external receipt.

123.2 Scientific dependency boundary

The emitter imports only the byte-pinned parser/metadata implementation from the older source machinery; it does not reuse its transition semantics. Test oracles independently use the frozen eager and sparse new-rule evaluators. The included prior proofs establish endpoint geometry, prospective-isolation involutivity, and preservation of admissible source simulation. The present report restates the rule and proves its direct-slot and arithmetic compilation, so reading it does not require an inaccessible earlier PDF.

The portable fixture revision changed only a provenance label. Its old/new digests and unchanged canonical non-provenance payload digest remain recorded. Any later test-output revision must similarly preserve original lineage and be checked by typed receipt comparison, rather than silently replacing evidence. Hashes establish version identity, not mathematical truth; the explicit arguments, independent code audit, and bounded computations have distinct roles.

124 Further research with distinct obligations

Three next directions should be kept separate. First, cheaper sorting, shared arithmetic subexpressions, special treatment of empty families, or more economical raw-slot routing may reduce the finite-horizon circuit. Each change must retain total deterministic assignments, canonical padding, and the same all-input raw/prospective semantics; it needs new exact ledgers and compiler-binding checks, not only faster evaluation. Second, guard-table compression might replace the explicit $(J + 2)^2$ scan with a smaller verified Boolean circuit. Its syntax, lookup or decoding, coefficients, and complete auxiliary fibers must be paid; a compressed table is not a free constant oracle.

Third, replacing the fixed horizon by an unknown or unbounded time parameter while keeping arity fixed is a different Diophantine problem. This report supplies neither a canonical bounded-arity history encoding nor a finite-fold proof for such an encoding. An exact-target or halting relation with existential time also needs its acceptance semantics specified. These questions remain unpaid mathematical obligations rather than consequences of the current quartic degree bound.

125 Conclusion

A concrete paid compiler now binds the new parallel particle rule to an ordinary quartic exact-endpoint relation for every fixed validated finite source, mass, horizon, and direction. Its unique complete natural witnesses follow from deterministic total gadgets and fully defined inactive computation. Direct source slots, all-type isolation, and full prospective rediscovery distinguish

the construction from an ordered-factor certificate. Exact resource records and the expanded pair example make the cost inspectable, while the five-particle separator supplies a meaningful semantic rejection test.

The remaining boundary is explicit: fixed-horizon arity grows with time, larger expanded polynomials and a universal-source instantiation have not been emitted, and mass-five horizons two and three have not been completed as certificate tests. None of these limitations weakens the stated finite-horizon theorem; none should be hidden when considering the construction as a Diophantine substrate.

126 File map for the scientific argument

[write] This section is Report 28’s Appendix A; the paths are delivery paths, mapped to shipped names in the README.

Relative item	Role
<code>research/PROOF.md</code>	All-input circuit binding, natural uniqueness, cost theorem
<code>research/compiler.py</code>	Direct raw slots, sorting, lanes, prospective tests, transport and composition
<code>research/circuit.py</code>	Integer/Boolean/comparison gates and polynomial expansion
<code>research/audit-design.md</code>	Independent geometric design analysis; old isolation macro cost superseded
<code>research/audit-emitter.md</code>	Independent concrete implementation audit
<code>research/resource-ledger.json</code>	Exact costs, coefficient magnitudes and serialized example bytes
<code>research/fixtures.json</code>	Full mandatory oracle snapshots and finite domains
<code>research/verify_example.py</code>	Independent ordinary-polynomial expansion and assignment check
<code>research/references/</code>	Pinned source schema, rule, compiler, lemma and preservation dependencies
<code>bitbound/</code>	Explicit complete-witness height proof, independent audit, and mandatory bounded tests
<code>research/PROVENANCE.json</code>	Byte-identity source and proof dependency records
<code>research/MANIFEST.json</code>	Scientific-packet inventory and source identities
<code>RELEASE-MANIFEST.json</code>	Exact release inventory and source lineage
<code>replay_release.py</code>	Normal/optimized external-output replay and source immutability checks

The exact release inventory, rather than this thematic map, determines the required files. Neither this report nor its archive contains a claim that historical ordered-rule tests certify arbitrary new-rule supports.

Part VII

Smaller sufficient radii for the parallel rule: two-scale recognition and shorter exactness windows (Reports 70 and 71)

[write] (Batch 91, 4 October 2026.) The two Reports of this Part keep the literal endpoint templates, the five-particle encoding and the whole admissible doubled micrograph of Report 26 (Part VI) and lower the sufficient radius of a two-involution full-shift extension of the same admissible dynamics. Report 70 separates a short geometric recognition radius a from the contextual control radius r in the edge block (isolation distance $\max\{2(b+a), b+r\}$ in place of Report 26's $2(b+r)$) and keeps Report 26's phase block: sufficient radius $108D + 149 + 3J$ in place of $180D + 258 + 9J$. Report 71 shortens triple exactness from $3b + 1$ to b in both blocks: $76D + 105 + 3J$, or $76D + 104 + 3J$ by a support refinement for the same rule. For the inherited universal ledger ($D = 509\,508$, $J = 0$) the bounds are 91 711 698 (Report 26), 55 027 013 (Report 70) and 38 722 713, or 38 722 712 (Report 71). Together they answer Report 26's first question (dated note in Section 100). Each Report defines a *new* full-shift rule: Report 26's F_{new} (written F_{26} in this Part), Report 70's F_* and Report 71's F_s agree on every admissible state and differ on explicit malformed inputs (Sections 135 and 145). Consequently Report 27's evaluator (Theorem 101.1) and Report 28's quartic certificates (Theorem 111.1), which are for Report 26's rule, neither evaluate nor certify F_* or F_s , and neither Report of this Part supplies an evaluator or a certificate; admissible semantic statements transfer through the proved agreement. The bounds are sufficient radii, not measured or minimum radii, and spatial radii, not arithmetic gate counts. Report 70 also classifies the injective number-conserving binary rules of radius two (Section 136), independently of its main proof.

The Hilbert's-tenth programme read both Reports in a scoped intake, `Computability/HilbertTenthProblem/Papers/research-wip/native-stream-queue/review_new_arithmetic_0d7f51c44.md` (commit `a9ab9a698`, 4 October 2026, together with six other archives of the same arrival). It read both delivery READMEs completely and, of the articles, Report 70's abstract and Sections 127, 134 and 137 and Report 71's abstract and Sections 140, 144 and the conclusion; it found no correction. It records that the bounds are sufficient radii, not arithmetic gates or minimum-radius theorems; that both Reports keep the absent universal-source table as an inherited premise; that earlier arbitrary-input transition certificates, circuit counts, ordered-factor budgets and evaluators do not transfer; and that Report 71's support-bound argument was read but neither Report 71's complete recognition and geometry proof nor Report 70's full-shift theorem was re-audited. It found in these archives no paid arithmetic construction below the programme's 84-operation unbounded ordinary-input bound, and notes that the radius results would first need a new arbitrary-input arithmetic implementation and cost audit. No other review of either Report exists; the audits and reviews shipped with them were made by the same AI-assisted pipeline. Claims of the two Reports, of their reviews and of this write that are not proved here are collected, with what is missing, in Section 149 at the end of this Part.

Two scale recognition for the binary five particle compiler

A smaller sufficient radius with exact trajectory preservation

Research report 70

4 October 2026

Abstract of Report 70, as delivered. We separate a short geometric recognition radius from a longer contextual-control radius in a parallel conservative involution. If local involutions write radius b , recognition reads radius a , and control reads radius r , with $b \leq a \leq r$, isolation at distance $H = \max\{2(b+a), b+r\}$ and a finite prospective recognition test give a full-shift involution of radius at most $b+a+H$. The proof preserves all recognition and selection decisions, including initially rejected anchors and cooperative candidate births.

Applying this lemma to the literal binary five-particle compiler, while retaining its phase block, gives sufficient radius $108D + 149 + 3J$ in place of Report 26's $180D + 258 + 9J$. The encoding, every admissible forward and inverse trajectory, exact clocks, reflections, halt observer and clean-target construction are unchanged. The inherited universal-source ledger yields 55 027 013 instead of 91 711 698, approximately a forty percent reduction in the proved upper bound. This is a new full-shift rule; an explicit four-particle input distinguishes it from Report 26. A separate complete finite certificate leaves only the five shifts among the 428 conservative binary five-input tables. Neither result is a priority or minimum-radius claim. No new-rule simulator, universal template array, arithmetic evaluator or universal-source execution is supplied.

[write] Delivered as `Two_scale_recognition_radius_sources.zip` (batch 91, manuscript 22; arrival 0d7f51c44), dated 4 October 2026; files placed by 22a8ca89e with the prefix `70-two-scale-`. No author is named and no ProveIt commit is pinned. Report 70 continues Report 26, whose archive it pins by SHA-256 (20a23b1ee22a..., the committed `Two_Parallel_Conservative_Involutions_Package.zip` of db37d18c8, whose files are this report's `26-parallel-involutions-` files). Its Theorem 128.1 refines Report 26's Lemma 90.1: for $a = r$ its constants are that lemma's, $H = 2(b+r)$ and radius $3(b+r)$, but the selection is by recognized anchor through a state-dependent local involution, not by typed key. Lemma 129.1, the guard-free anchor argument of Section 132, the witness of Section 135 and Theorem 136.1 are new here; Lemma 131.1 restates Report 26's phase block, and Sections 130, 131 and 133 restate the source interface, templates, clocks and clean targets of Reports 15 and 26 (notes there). Report 70 writes α for the number of zero-update branches, which Parts I–VI call a , because it reserves a for the recognition radius (Section 2). Its macros `\N` (printing $\mathbb{Z}_{\geq 0}$) and `\Fold`, `\Eold`, `\Pold` (printing F_{26} , E_{26} , P_{26}) are `\Nge`, `\Ftwosix`, `\Etwosix`, `\Ptwosix` here, because `\Fold` prints F_{old} in Part VI; nothing printed changed. Its table of contents is omitted. Its four architecture references carry the same locators as Report 26's and are the existing entries; its Morita reference is the existing entry [4], and entries for Fuks–Sullivan and Schranko–de Oliveira are added.

127 Objects and principal result

A configuration is $x \in \{0, 1\}^{\mathbb{Z}}$. A particle is a site whose bit is one. A cellular automaton (CA) is a translation-equivariant map with a finite, uniform local radius. Number conservation means preservation of the ordinary number of ones on every finite-particle configuration. We do not equate divergent sums on infinite inputs.

The starting point is Report 26's two-block rule $F_{26} = P_{26} \circ E_{26}$. Its finite endpoint templates come from the earlier guarded-swap compiler. We retain those literals, their type-specific write windows, their guards and their signed five-particle encoding. We change only the edge block's full-shift selection mechanism, and retain P_{26} exactly.

Theorem 127.1 (Source uniform radius reduction). *For every finite source satisfying the interface in Section 130, there is a binary CA*

$$F_* = P_{26} \circ E_*$$

with these properties.

1. Both E_* and P_{26} are involutions on the entire binary full shift, and each preserves ordinary finite particle number.
2. The forward and inverse maps agree with Report 26 on the entire admissible doubled micrograph: all natural-counter homes, both signs of all strict intermediates of enabled source edges, and all translations.
3. If m is the number of controls, p the number of nonzero-update branches, and J the common finite counter-class cut, then

$$R(F_*) \leq 108D + 149 + 3J, \quad D = 2m + 4p. \quad (127.1)$$

4. The inverse is $E_* \circ P_{26}$. Each of the two individually conservative involutions is a time reverser for F_* .

The geometric recognition predicate ignores contextual guards but retains each template's exactness window. Different types at the same integer anchor contribute *one* recognized anchor. A symmetric before-and-after uniqueness test makes the guarded choice at that anchor reversible even on malformed inputs. It is essential to prove geometric anchor uniqueness independently of guarded type uniqueness; Sections 131–132 do so directly from every literal endpoint.

The main theorem is source-uniform and its proof is mathematical. The numerical universal corollary inherits exactly Report 26's source and normalization dependencies; this report does not re-establish or materialize that universal source. Section 134 states the complete ledger and dependency boundary. Section 135 shows that the new rule and Report 26 really differ off the admissible set. The independent radius-two classification in Section 136 does not enter the radius-reduction proof.

128 The two scale full shift lemma

128.1 Recognition and reversible local control

Write $I_t(u) = [u - t, u + t] \cap \mathbb{Z}$. Fix integers $0 \leq b \leq a \leq r$ and a translation-covariant Boolean predicate $\rho(x, u)$ determined by $x|_{I_a(u)}$. Its recognized-anchor set is

$$\mathcal{R}(x) = \{u \in \mathbb{Z} : \rho(x, u) = 1\}.$$

For each integer u , let T_u be the corresponding translate of a fixed everywhere-defined configuration involution. Assume:

1. T_u fixes every coordinate outside $I_b(u)$
2. Its replacement word on $I_b(u)$ is determined by $x|_{I_r(u)}$
3. It preserves the sum of bits on $I_b(u)$
4. $T_u x \neq x$ implies $\rho(x, u) = 1$

The map T_u may be the identity on an input, including a recognized input. It need not be a fixed transposition chosen independently of x ; its entire state-dependent definition must be an involution.

Set

$$H = \max\{2(b + a), b + r\}. \quad (128.1)$$

Select $u \in \mathcal{S}(x)$ precisely when

$$u \in \mathcal{R}(x), \quad (128.2)$$

$$\mathcal{R}(x) \cap I_H(u) = \{u\}, \quad (128.3)$$

$$\mathcal{R}(T_u x) = \mathcal{R}(x). \quad (128.4)$$

Define $A(x)$ by applying all selected local replacements simultaneously and leaving every other site unchanged. The last test compares recognition before and after just one hypothetical T_u . It is neither a recursive eligibility test nor a comparison restricted to currently selected anchors.

Theorem 128.1 (Separated recognition and control). *The map A is a translation-equivariant, finite-number-conserving involution on $\{0, 1\}^{\mathbb{Z}}$. For every bi-infinite configuration,*

$$\mathcal{R}(Ax) = \mathcal{R}(x), \quad \mathcal{S}(Ax) = \mathcal{S}(x), \quad A^2 x = x.$$

Its radius is at most

$$b + a + \max\{2(b + a), b + r\}. \quad (128.5)$$

128.2 Finite local meaning of the prospective test

A change in $I_b(u)$ can affect $\rho(x, v)$ only when its read interval meets that write interval, hence only for $|v - u| \leq b + a$. Thus (128.4) is equivalent to the finite family

$$\rho(T_u x, v) = \rho(x, v) \quad (v \in I_{b+a}(u)).$$

These recognition windows have union $I_{b+2a}(u)$. Constructing the hypothetical replacement additionally reads $I_r(u)$, so the prospective test is determined by

$$I_s(u), \quad s = \max\{r, b + 2a\}. \quad (128.6)$$

Selected anchors have pairwise distance greater than $H \geq 2b$. Their write intervals are therefore disjoint. This defines A coordinatewise even for an arbitrary dense infinite input.




Figure 8: Read and write roles at an isolated anchor. The schematic shows the compiler regime $r \geq b + 2a$; the abstract theorem allows either order. The separation conditions are $H \geq b + r$ and $H \geq 2b + 2a$. Isolation itself reads $I_{H+a}(u)$.

128.3 Recognition invariance on arbitrary inputs

Fix any potential recognition anchor v , whether initially recognized or absent. Its read interval $I_a(v)$ meets at most one selected write interval: meeting those at distinct u, w would imply

$$|u - w| \leq |u - v| + |v - w| \leq 2(b + a) \leq H,$$

contradicting isolation. After the simultaneous update, the word read by $\rho(\cdot, v)$ is therefore its original word or exactly its word after a single selected T_u . In the latter case the prospective test at u preserves its value. Consequently $\mathcal{R}(Ax) = \mathcal{R}(x)$ on the full shift. This explicitly excludes cooperative births or deaths caused jointly by different selected replacements.

128.4 Every selection decision is invariant

Consider *any* recognized isolated anchor u , including one whose prospective test fails. Every selected $w \neq u$ is another recognized anchor and hence satisfies $|w - u| > H$. Since

$$H \geq b + r, \quad H \geq 2b + 2a,$$

its write interval misses both $I_r(u)$ and $I_{b+2a}(u)$. Therefore every other selected update leaves the full local-control and prospective-comparison data at u unchanged.

If u is unselected, no update at u occurs either. Its prospective test remains false, so it stays unselected. If u is selected, the test data after all updates are those in $T_u x$. Because T_u is an involution, testing there interchanges the two sides of

$$\mathcal{R}(T_u x) = \mathcal{R}(x).$$

The equality remains true. Recognition invariance preserves isolation. Nonisolated recognized anchors remain nonisolated, and absent anchors remain absent. Hence $\mathcal{S}(Ax) = \mathcal{S}(x)$, including every rejected status.

A recognized isolated anchor with $T_u x = x$ passes prospectivity and is selected, but writes nothing. Its local identity choice remains unchanged because other selected writes miss $I_r(u)$. Thus identity actions are covered rather than tacitly removed from the selected set. The proof does not assert preservation of every internal guarded-template label; it preserves exactly the recognition and selection data needed for inversion.

128.5 Involution conservation and radius

At the second application, the same anchors are selected. The local-control word at each selected u contains its own first image and no modification from any other selected write. The involution T_u therefore restores that write block. All other sites remain fixed, proving $A^2 = \text{id}$.

Disjoint equal-weight finite replacements conserve ordinary finite mass. In particular, on a finite input a nonidentity replacement requires at least one input particle in its block; only finitely many such disjoint blocks can act. Translation covariance is inherited from the predicates and translated local maps.

For an output coordinate i , only anchors in $I_b(i)$ can write it. Isolation at such an anchor reads radius $H + a$, local control reads r , and prospectivity reads $\max\{r, b + 2a\}$. But

$$H + a \geq r, \quad H + a \geq b + 2a.$$

Thus the output depends only on $I_{b+H+a}(i)$, proving (128.5). This finishes Theorem 128.1. No finite-support reduction, source semantics, empirical trajectory argument or proof-assistant execution is used.

129 A reversible choice among templates

Fix a finite type family G . Type g has a write set $W_g \subseteq I_b(0)$ and two distinct equal-weight endpoint words. Let $\tau_{g,u}$ transpose these words on $u + W_g$, fix every other word there, and fix all coordinates outside that type-specific window. This is an everywhere-defined involution.

Let $c_g(x, u)$ be its exact guarded candidate predicate, reading at most $I_r(u)$ and implying that the endpoint word is present. At one integer anchor put

$$C_u(x) = \{g \in G : c_g(x, u) = 1\}.$$

Define

$$T_u x = \begin{cases} \tau_{g,u} x, & C_u(x) = \{g\} \text{ and } C_u(\tau_{g,u} x) = \{g\}, \\ x, & \text{otherwise.} \end{cases} \quad (129.1)$$

Both uniqueness tests range over *all* edge types at the same anchor. There is no priority rule or tie-breaking choice.

Lemma 129.1 (Mutual uniqueness). *The map in (129.1) is an everywhere-defined, equal-weight involution. It writes only $I_b(u)$ and its replacement word is determined by $I_r(u)$.*

Proof. If the rule acts with type g , its image y has $C_u(y) = \{g\}$ and $\tau_{g,u} y = x$, with $C_u(x) = \{g\}$. The same symmetric test therefore acts in reverse. If the rule does not act, it fixes its input and fixes it again. This covers zero candidates, competing candidates and candidate creation in the hypothetical image. Each possible action preserves weight. Evaluating $c_h(\tau_{g,u} x, u)$ needs no expanded radius: inside W_g the replacement is known, and every unchanged bit is already in $I_r(u)$ because $b \leq r$. Thus the before-and-after choice uses only that original radius- r word. $\square$

This construction does not require source-domain or source-image disjointness to be globally reversible. Those hypotheses enter later, solely to prove that the intended admissible transition is accepted. For an empty family, take every C_u and $\mathcal{R}$ empty and every $T_u = \text{id}$; all uniform upper bounds still apply.

130 The source and five particle geometry

130.1 The complete source interface

The source is a finite deterministic partial injection on $Q \times \mathbb{Z}_{\geq 0}^2$. Its m controls include a designated halt control h with no outgoing branch. A branch $e : q \rightarrow q'$ selects the left or right counter, with side $v \in \{-1, +1\}$ and index j , changes it by $\Delta \in \{-1, 0, +1\}$, and has a Boolean domain guard G_e . Guards use finitely many tests $c_j = k$ or $c_j > k$. Enabled updates have natural outputs. Branch domains at each source control are disjoint, and *actual branch images* at each target control are disjoint, on all natural-counter IDs, including unreachable ones.

The exact image predicate is

$$I_e(c') = G_e(c' - \Delta \mathbf{e}_j) \wedge (c' - \Delta \mathbf{e}_j \in \mathbb{Z}_{\geq 0}^2). \quad (130.1)$$

Fix $J \geq 0$ such that every domain and every exact image predicate is determined by the classes $0, 1, \dots, J, > J$ of both counters. Write p for the number of nonzero-update branches and α for the number of zero-update branches. We reserve a for the recognition radius. Other controls may be stuck on some IDs; stuckness at a nonfinal control is not designated halting.

These are the authenticated compiler's mathematical source assumptions. In particular, an exact reverse guard cannot be replaced by the forward guard. The construction below specifies finite Boolean predicates directly; no source parser or compiler program is executed to justify them.

130.2 Signed modes and encoding

There are $m + 2p$ unsigned modes: H_q for each control, and O_e, I_e for each nonzero branch. Give each signed mode a distinct gap $d(u^\pm) \in \{1, \dots, D\}$, where

$$\begin{aligned} D &= 2m + 4p, & S &= 2D + 2, & B_2 &= D + 1, & L &= 3D + 4, \\ b &= B_3 = L + D + 1 = 4D + 5, & K &= 3b + 1, & Z &= 10b + 10 + 2J. \end{aligned} \quad (130.2)$$

A head at anchor x consists of $\{x, x + d(u^\pm)\}$. Its anchor is always the leftmost head particle, also during leftward motion. At origin zero the home encoding is

$$\text{Enc}(q, c_\ell, c_r)^\pm = \{-(Z + c_\ell), 0, Z + c_r, S, S + d(H_q^\pm)\}. \quad (130.3)$$

The three singletons are unlabelled left endpoint, origin and right endpoint. Their gaps are at least Z . The head is the unique pair at separation at most D ; all head-to-marker distances exceed D . The configuration has exactly five particles in a binary zero background. Neither an extra track nor a spatial parity pattern, external phase schedule or clock particle is present.

130.3 The precise admissible graph

An enabled zero-update branch is one unsigned home-to-home microedge. For a nonzero enabled branch, with origin o and selected old marker m , set $m' = m + v\Delta$. Its subdivision is:

1. Dispatch the home head from $o + S$ to mode O_e at $o + vS$, using G_e on old counters
2. Move the outbound anchor by v until it reaches $m - vS$
3. Replace marker m by m' and change to I_e at $m' - vS$
4. Move the inbound anchor by $-v$ until it reaches $o + vS$
5. Commit to $H_{q'}$ at $o + S$, using I_e on post counters

The unsigned admissible set contains all homes in $Q \times \mathbb{Z}_{\geq 0}^2$ and *only* the strict intermediate configurations of enabled subdivisions. A merely geometric moving packet with a false branch guard is excluded. The doubled set contains both signs of every such node and all translations.

Intermediate nodes retain the branch label. Outbound modes retain old counters, inbound modes retain post counters, and the head anchor locates the position on a unique corridor. Distinct branches cannot share an intermediate node. Exact source-image disjointness prevents merging at homes. The unsigned micrograph is therefore a partial injection. For every unsigned edge $c \rightarrow c'$, its doubled edge joins c^+ to c'^- ; the phase matching joins the two signs of each node.

131 All literal templates and local predicates

Every displayed support below is a binary word with exactly those occupied sites and zeros elsewhere in its *type-specific* interval $W_g = I_{B_g}(0)$. Pair types have $B_g = B_2$; triple types have $B_g = b$. The invariant template anchor is zero in the display, not necessarily a particle or the current head anchor. The map $\tau_{g,u}$ exchanges just these words on $u + W_g$ and retains all other coordinates.

131.1 Edge templates

For each moving mode u , its forward direction is $w = v$ for O_e and $w = -v$ for I_e . Its free pair is

$$\{0, d(u^+)\} \longleftrightarrow \{w, w + d(u^-)\}. \quad (131.1)$$

The invariant anchor is the *predecessor* head anchor. Marker-contact travel uses

$$\{0, wt, wt + d(u^+)\} \longleftrightarrow \{0, w(t+1), w(t+1) + d(u^-)\}, \quad (131.2)$$

$$S \leq t \leq L,$$

$$\{0, -wt, -wt + d(u^+)\} \longleftrightarrow \{0, -w(t-1), -w(t-1) + d(u^-)\}, \quad (131.3)$$

$$S+1 \leq t \leq L.$$

The first family is behind travel and the second ahead travel; both are anchored at the contacted marker. For a nonzero branch there are three interactions:

$$\{0, S, S + d(H_q^+)\} \longleftrightarrow \{0, vS, vS + d(O_e^-)\}, \quad (131.4)$$

$$\{0, -vS, -vS + d(O_e^+)\} \longleftrightarrow \{v\Delta, v\Delta - vS, v\Delta - vS + d(I_e^-)\}, \quad (131.5)$$

$$\{0, vS, vS + d(I_e^+)\} \longleftrightarrow \{0, S, S + d(H_{q'}^-)\}. \quad (131.6)$$

Dispatch (131.4) uses G_e , the endpoint (131.5) is unguarded, and commit (131.6) uses I_e . The endpoint anchor remains the *old* marker coordinate even when the actual marker moves. Each zero-update branch contributes

$$\{0, S, S + d(H_q^+)\} \longleftrightarrow \{0, S, S + d(H_{q'}^-)\}, \quad (131.7)$$

guarded by $G_e = I_e$.

Every support fits in its specified write interval. Every triple has exactly one close pair of gap at most D , and its singleton has distance greater than D from both pair particles. This remains true at distance $L+1$ and at shifted endpoint images. Distinct signed gaps make the two endpoints distinct and nontranslates.

131.2 Raw recognition and guarded control

For edge type g , put $\ell_g = 3B_g + 1$. Define

$$\rho_g(x, u) = 1 \iff \{i \in I_{\ell_g}(u) : x_i = 1\} \in \{u + P_g, u + Q_g\}. \quad (131.8)$$

The short predicate and guarded type set are

$$\rho(x, u) = \bigvee_{g \in G_E} \rho_g(x, u), \quad C_u(x) = \{g : \rho_g(x, u) \text{ and its contextual guard holds}\}.$$

Recognition has radius $a = K = 3b + 1$. It retains the *individual* pair exactness radius L and triple exactness radius K . Taking every pair to be exact on the common larger write interval I_b would be wrong: a valid free head can have a marker at distance $L+1 < b$. Such a particle must survive a pair replacement.

Home-interaction guards read the sites

$$u - (Z + k), \quad u + (Z + k), \quad 0 \leq k \leq J.$$

For either counter, exactly one occupied read selects its finite class; all-zero reads mean $> J$; multiple occupied reads make the guard false. These deterministic finite-word conventions specify guards even on malformed inputs. They read outside every write interval and remain unchanged by an endpoint swap. All exactness and guard predicates read at most

$$r = Z + J = 10b + 10 + 3J.$$

Use (129.1) to define T_u . It is a radius- r controlled involution with write bound b . A nonidentity action has a guarded candidate and hence a raw recognized anchor. Theorem 128.1 therefore applies and defines E_* on every bi-infinite binary word.

131.3 The unchanged phase block

For each moving mode retain the phase pair

$$\{0, d(u^+)\} \longleftrightarrow \{0, d(u^-)\},$$

and, for $s \in \{-1, +1\}$ and $S \leq t \leq L$, the phase triples

$$\{0, st, st + d(u^+)\} \longleftrightarrow \{0, st, st + d(u^-)\}.$$

For each control retain the home-phase triple

$$\{0, S, S + d(H_q^+)\} \longleftrightarrow \{0, S, S + d(H_q^-)\}.$$

There are no contextual phase guards. To specify P_{26} *exactly*, keep Report 26's typed keys, rather than coalescing them into anchors. A phase key is $k = (g, u)$ satisfying the radius- ℓ_g exactness predicate. Let $\mathcal{C}_P(x)$ be the full typed-key set. Put $r_P = K$ and $H_P = 2(b + r_P)$. Select k when there is no other typed key with anchor within H_P of u and

$$\mathcal{C}_P(\tau_k x) = \mathcal{C}_P(x). \quad (131.9)$$

This equality is a finite check on keys within $b + r_P$ of u . Apply all selected transpositions simultaneously. This is the unchanged Report 26 definition.

Lemma 131.1 (Retained phase involution). *The map P_{26} is a full-shift conservative involution with radius at most $3(b + r_P) = 12b + 3$.*

Proof. Selected typed keys have distinct anchors more than $2(b + r_P)$ apart, hence disjoint write intervals. A typed candidate's radius- r_P read meets at most one selected write, so (131.9) preserves the entire candidate set, including absent keys. The prospective test at u reads only $I_{b+2r_P}(u)$. For any isolated typed key, every other selected write misses that interval. A rejected key's test is unchanged; a selected key's own transposition interchanges the two candidate sets in the equality. Recognition and isolation are invariant, as are all selection decisions. A second application reverses the same disjoint transpositions. Equal weight gives finite conservation. Isolation reads $H_P + r_P$ about its key, so an output coordinate reads at most $b + H_P + r_P = 3(b + r_P)$. $\square$

This typed-key proof applies to dense and infinite configurations and to an empty family. It avoids assuming that the earlier *ordered* phase-gate product was itself a full-shift involution. Only the already parallel Report 26 phase block is retained here.

[write] P_{26} is Report 26's phase block $P_{||}$ of Theorem 89.1 (Section 93), and Lemma 131.1 is Report 26's Lemma 90.1 applied with read radius $r_P = K = 3b + 1$; its radius $R_P = 12b + 3$ is Report 26's (Section 95). The proof is restated, not new.

132 Agreement on every admissible state

The global involution proof is independent of source semantics. To retain the intended computation, we must additionally prove that all intended local transitions pass the new filters. Guarded type uniqueness alone does not prove the required *guard-free anchor* statement. We establish it directly from the literal templates.

132.1 Every raw match uses the actual head

At every admissible node there is exactly one particle pair with gap at most D , the head. Markers are separated by at least Z ; head-marker distances are at least $S - D = D + 2 > D$. Every endpoint template has exactly one close pair, so every raw match must use that actual head. Its distinct gap determines the signed mode. A triple's remaining displayed particle must be one of the three markers.

A fixed triple has diameter at most $2b < Z$, but that observation alone would not exclude *different* types using different nearby markers. The literal formulas give the stronger fact that every triple places its singleton within $L + 1$ of the actual head anchor. Because $Z > 2(L + 1)$, at most one marker can lie in that range. The signed mode, side and distance then determine the possible travel type, subject to the partitions below. The third marker cannot supply a remote alternative.

Other markers lie outside the intended triple's exactness window. For a current-marker anchor their distances are at least $Z > K$. For a reverse endpoint's old-marker anchor they are at least $Z - 1 > K$, since the actual marker moved by at most one. All intended triple exactness conditions therefore hold. Free-pair exactness is addressed separately by its own radius L .

132.2 Homes and same anchor branch alternatives

At a plus home H_q^+ , only forward dispatch or direct types with source q have the correct signed gap. Their home endpoint has head anchor $u + S$ and singleton u , forcing the integer anchor to be the origin. At a minus home H_q^- , only reverse commit or direct types with target q match, again at that origin. No moving free pair, moving travel triple or other interaction has the same signed gap.

Omitting guards may thus expose several raw *types* at a home, but all have the same integer anchor. The raw recognized set is empty or the singleton origin. Its guarded set $C_o(x)$ is empty or a singleton: source-domain disjointness applies at plus homes, and exact-image disjointness at minus homes. The class reads recover the actual counter classes because the interaction head stays within $b < Z$ of the origin.

If reverse commit is enabled at post counters c' , predicate (130.1) supplies a genuine natural predecessor $c = c' - \Delta \mathbf{e}_j$ satisfying $G_e(c)$. It enters that enabled subdivision, not an arbitrary guard-invalid moving packet. Forward dispatch similarly uses the true domain. Direct zero-update transitions use $G_e = I_e$. A blocked home may have raw recognition while all guarded candidates are absent; then its local map is identity, which is harmless.

132.3 Both signed corridor partitions

Orient the moving corridor in its forward direction w . Let A be its departure marker, C its arrival marker, and x the head anchor. Put

$$N = w(C - A) \geq Z, \quad t = w(x - A), \quad S \leq t \leq N - S.$$

For outbound travel $w = v$ and $N = Z + c$. For inbound travel $w = -v$ and $N = Z + c + \Delta$. Natural-output guards ensure the latter is still at least Z . This notation covers both physical directions without reflecting the head-anchor convention.

For plus states, the free predecessor key is x , so pair exactness means $L < t < N - L$. The literal triple ranges complete the following disjoint exhaustive partition:

Plus case	Oriented head position	Raw anchor
Behind travel	$S \leq t \leq L$	A
Free travel	$L + 1 \leq t \leq N - L - 1$	x
Ahead travel	$N - L \leq t \leq N - S - 1$	C
Forward interaction	$t = N - S$	as below

For minus states, the free inverse key is $x - w$, with oriented coordinate $t - 1$. Hence pair exactness means $L < t - 1 < N - L$, yielding:

Minus case	Oriented head position	Raw anchor
Reverse interaction	$t = S$	as below
Behind inverse travel	$S + 1 \leq t \leq L + 1$	A
Free inverse travel	$L + 2 \leq t \leq N - L$	$x - w$
Ahead inverse travel	$N - L + 1 \leq t \leq N - S$	C

These intervals include every equality case and have no overlaps or gaps. A behind triple from plus $t = L$ arrives at minus $t = L + 1$, but its free predecessor key still sees marker A at distance L , so a free inverse cannot compete. An ahead free move begins at marker distance $L + 1$ and ends at minus distance L ; reverse ahead triples reach only distance $L - 1$. At the forward wall there is no travel beyond distance S , and at the reverse wall no travel from behind distance $S - 1$ is permitted.

Within a triple range, the signed gap, the unique possible marker, its side and its integer distance specify exactly one type. The free pair is the unique type on its free interval. All these statements concern the guard-free endpoint geometry.

132.4 Interaction walls and old marker anchors

The four moving walls have these forced interactions:

1. Outbound minus at $t = S$: reverse dispatch at the origin A
2. Outbound plus at $t = N - S$: forward endpoint at the old selected marker $m = C$
3. Inbound plus at $t = N - S$: forward commit at the origin C
4. Inbound minus at $t = S$: reverse endpoint at the old marker coordinate $m = m' - v\Delta$

The last anchor is not generally the actual current marker m' . In (131.5), its singleton has offset $v\Delta$ and its head anchor has offset $v\Delta - vS$. Both force the same old anchor in the reverse orientation.

For the matching mode, the candidate singleton location forces precisely the stated wall. Reverse outbound dispatch requires a marker at $x - vS$, forward outbound endpoint at $x + vS$, inbound commit at $x - vS$, and reverse inbound endpoint a current marker at $x + vS$. In the oriented corridor those are respectively the behind or ahead wall indicated above. The third marker cannot occupy a competing contact location. Branch-specific signed O_e/I_e gaps exclude all other branch interactions. The corridor tables exclude travel at the same wall.

Reverse dispatch retains the old counters of its enabled source edge, so G_e is true there. Forward commit retains that edge's post counters, so its exact I_e is true. Endpoint interactions need no contextual guard. Thus every admissible moving node has exactly one raw anchor and exactly its intended guarded type there. The definition of admissibility, restricting intermediates to *enabled* subdivisions, is essential to this guard argument.

132.5 The two endpoints keep the same recognized anchor

Every intended microedge has two admissible signed endpoints. Their invariant raw anchor is:

- The predecessor head anchor for free travel
- The contacted unchanged marker for triple travel
- The unchanged origin for dispatch, commit and direct transitions
- The old endpoint-marker coordinate for endpoint transitions in both orientations

At an intended edge of type g and anchor u , the previous arguments give

$$\mathcal{R}(x) = \{u\} = \mathcal{R}(\tau_{g,ux}), \quad C_u(x) = \{g\} = C_u(\tau_{g,ux}).$$

At homes several raw types can coexist, but they still count as just this one anchor. The local mutual-uniqueness test therefore chooses the intended transposition. Isolation and prospective raw-anchor invariance pass automatically. At a missing edge the local maps are identity whether or not raw recognition remains. Hence E_* equals the intended doubled edge matching on *every* admissible input.

For the retained phase block, a home has its unique home-phase triple. A moving state has the free phase pair when every marker is farther than L from its unchanged head anchor, and otherwise the unique phase triple at distance $S, \dots, L$ from the one nearby marker. The phase action fixes all markers, the head anchor and unsigned mode. Thus its free/near/home case and typed key are unchanged, and its filters pass. It is the sign flip throughout the admissible set.

Both restricted blocks preserve that set. Report 26 has exactly the same doubled edge matching and phase flip there. Consequently $F_* = F_{26}$ and $F_*^{-1} = F_{26}^{-1}$ on the entire admissible doubled micrograph, proving the agreement assertion of Theorem 127.1. This is not a claim about only one forward orbit or a bounded set of counters.

133 Exact clocks observations and targets

On an admissible plus node with a successor, F_* advances one unsigned microedge and retains plus. On a minus node with a predecessor, it goes back one microedge and retains minus. At a missing edge it flips sign in one step. The two mathematical submaps are composed into one time-independent CA step.

For a nonzero branch at old selected counter c , the outbound and return travel lengths are $Z + c - 2S$ and $Z + c + \Delta - 2S$. Including the three interactions gives the unchanged clock

$$\tau_e(c) = 3 + 2(Z + c) + \Delta - 4S. \quad (133.1)$$

A zero-update branch takes one microedge. The proof preserves each physical step, not just source-level reachability or an asymptotic slowdown.

Since both factors are full-shift involutions,

$$F_*^{-1} = E_* P_{26}, \quad E_* F_* E_* = F_*^{-1}, \quad P_{26} F_* P_{26} = F_*^{-1}.$$

For a predecessor-free start and a finite forward path of T microedges, the reflected cycle has exactly $2T + 2$ states. Partial injectivity excludes any earlier repeated unsigned node: a repetition would propagate backward to a predecessor of the start. An infinite forward path has no repetition and no positive return. Nonfinal stuck homes also cause reflection. Consequently positive return

is equivalent to designated halting only under the additional hypothesis that nonhalting inputs have infinite forward source paths, with no nonfinal blocking.

The anchored halt word on $[0, S + D]$ has ones precisely at $0, S, S + d(H_h^+)$. Its length is

$$3D + 3.$$

The counter endpoints lie outside this word. Unique signed gaps distinguish it from every other mode. On valid trajectories it occurs exactly when the forward source run reaches h : a backward traversal could not first enter a new halt home without an outgoing edge from h , which has no exits.

For a clean exact target, require a *control-wide* no-incoming initial control q_0 , and a no-outgoing halt h . A single predecessor-free input ID is insufficient for the wrapper. Use controls $F(q), B(q)$ and a terminal H . Include each branch forward, its inverse backward with domain I_e , and the zero-update true-guard bridges

$$F(h) \longrightarrow B(h), \quad B(q_0) \longrightarrow H.$$

Reverse branches exchange domains and exact images. The control-wide empty boundary slots make the bridges compatible with determinism and partial injectivity on all IDs. The new counts are

$$m' = 2m + 1, \quad p' = 2p, \quad \alpha' = 2\alpha + 2,$$

with the same J . Recompute *all* geometry constants for this enlarged source. The target $\text{Enc}(H, c_\ell, c_r)^+$ restores the input counters; it is input-dependent, not one constant universal target. If the original run reaches h , the first bridge starts its exact reversed history. The no-incoming control q_0 cannot occur internally, so the second bridge fires exactly after that history has been undone. If the original run never reaches h , neither bridge can lead to H . Later CA reflections only retrace the same component. Thus target occurrence is equivalent to original designated halting.

For a halting run, let Θ be its original forward physical microedge count using those enlarged constants. Equation (133.1) satisfies $\tau(c, \Delta) = \tau(c + \Delta, -\Delta)$, and zero-update clocks are one in either direction. Thus the first target time is $2\Theta + 2$, and the full reflected period is $4\Theta + 6$. For the inherited one-increment clean example, $D = 18$, $J = 1$, $S = 38$, $Z = 782$ and $\Theta = 1416$, giving times 2834 and 5670. These values are semantic consequences of the unchanged admissible paths; no trajectory has been rerun here.

[write] Equation (133.1) is Report 15's clock (5) and Report 26's (94.1); the clean-target wrapper is Report 15's (Section 13), and its numbers are those of Report 26's clean-target example (Section 94).

134 The exact radius and resource ledger

For the edge block, substitute

$$a = 3b + 1, \quad r = Z + J = 10b + 10 + 3J.$$

The difference

$$(b + r) - 2(b + a) = 3b + 8 + 3J \geq 0$$

gives $H = b + r$. Therefore

$$\begin{aligned} R(E_*) &\leq b + a + H = 5b + Z + J + 1 = 15b + 11 + 3J, \\ R(P_{26}) &\leq 12b + 3, \\ R(F_*) &\leq 27b + 14 + 3J = 108D + 149 + 3J. \end{aligned} \tag{134.1}$$

This proves the radius assertion of Theorem 127.1. Report 26 instead used edge typed-key recognition with the common read radius r , exclusion $2(b + r)$ and edge radius $3(b + r)$. Its total was $180D + 258 + 9J$. The present decrease is

$$72D + 109 + 6J.$$

These are comparisons of *sufficient upper bounds*, not measured local radii or a proof of optimality. The edge safety mechanism changes the full-shift rule.

Quantity	Increment source	Clean sample	Universal ledger
m, p, α	2, 1, 0	5, 2, 2	122 622, 66 066, 75 495
D, J	8, 0	18, 1	509 508, 0
b	37	77	2 038 037
$a = 3b + 1$	112	232	6 114 112
$r = Z + J$	380	783	20 380 380
$H = b + r$	417	860	22 418 417
New edge bound	566	1 169	30 570 566
Unchanged phase bound	447	927	24 456 447
New total bound	1 013	2 096	55 027 013
Report 26 total bound	1 698	3 507	91 711 698
Decrease	685	1 411	36 684 685
Endpoint template count	95	353	269 291 358 255
Observer length	27	57	1 528 527

Table 15: Exact integer substitutions in source-uniform formulas. The last column inherits the same universal-source premise as Report 26 and does not assert a newly materialized universal rule. Each valid encoding still has exactly five particles over $\{0, 1\}$.

There are still

$$8pD + 29p + m + \alpha \quad (134.2)$$

endpoint templates across the two families. Indeed there are $4p$ pair types, $8pD + 23p + m$ unguarded triples and $2p + \alpha$ contextual triples. Each of $2p$ moving modes contributes behind positions $S, \dots, L$, ahead positions $S + 1, \dots, L$, and two sides of phase positions $S, \dots, L$; $L - S = D + 2$ yields the formula. These are alternatives evaluated within two global involutions, not two inexpensive arithmetic gates. Boolean work may still depend on the type count.

134.1 The inherited universal numerical corollary

The unchanged Report 19 receipt records

$$m = 122\,622, \quad p = 66\,066, \quad \alpha = 75\,495, \quad J = 0,$$

so $D = 509\,508$. Substitution gives 55 027 013, a decrease of 36 684 685 from Report 26's 91 711 698, approximately forty percent. The cited source table has SHA-256

fa61d06178d718511232a80a05e02192d940fde3c273f0c202f0623253b2ade3.

That table is *not* contained in this packet or the recovered Report 26 archive. Its authenticated receipt, counts and prior audit are retained. The source's historical construction, universality and required normalizations remain inherited premises, exactly as in Report 26. No universal source, interpreter, template array, truth table or new-rule evaluator was constructed or run for this report.

[write] Report 70's copy of that receipt is byte-identical to `data/19-lazy-eval-evidence-universal-receipt.json` (Part IV) and is not shipped again. The table `fa61d06178d7...` is

Report 17's `source.json`: absent from Report 70's package, as Report 70 says, but rebuilt byte for byte from shipped files of this report (README, "Reconstructing the excluded data").

The source-uniform radius proof does not need a universal table. Conversely, the arithmetic substitution does not supply a new universality proof. The finite-pattern and clean-target consequences transfer through the inherited source theorem and the proved admissible agreement, with their original hypotheses. The stronger periodicity/return consequence retains the nonfinal-blocking condition.

135 A four particle input separates the full shift rules

Take ordered controls (q, h) , one true-guard right increment $q \rightarrow h$, and $J = 0$. Its exact image guard is right-counter positivity, expressible by classes $0, > 0$. The literal inherited signed-mode order assigns gaps

$$H_q^+ : 1, H_q^- : 2, H_h^+ : 3, H_h^- : 4, O_e^+ : 5, O_e^- : 6, I_e^+ : 7, I_e^- : 8.$$

Here $D = 8$, $b = 37$, $L = 28$, $a = 112$ and $Z = r = 380$. Let

$$X = \{0, 5, 500, 505\}.$$

Each cluster has just two particles and the intercluster gap is at least $495 > 2b$. No triple can fit. The only close pairs have gap five, so the only edge raw anchors and guarded candidates are free O_e at 0 and 500. Their endpoint words are

$$\{0, 5\} \longleftrightarrow \{1, 7\}$$

and the 500-translate. Their own radius- L exactness windows contain precisely their pair.

Report 26's edge exclusion is $2(b + r) = 834$; the two candidates reject one another. Thus $E_{26}(X) = X$. The new exclusion is $H = b + r = 417$, so both anchors are isolated. At each anchor the guarded type is unique before and after its hypothetical swap. Its reverse free key remains the same predecessor key: the new leftmost particle is $u + 1$, so subtracting its direction one recovers u . No cross-cluster close pair or triple can arise. Both individual hypothetical swaps preserve the recognized set $\{0, 500\}$. Hence

$$E_*(X) = \{1, 7, 501, 507\}.$$

The unchanged phase exclusion is $H_P = 2(b + 3b + 1) = 298$. On X the free phase anchors are 0, 500; on $E_*(X)$ they are 1, 501. Their separation is 500, and each hypothetical sign flip keeps the same phase typed-key set. The phase pair exchange is $\{0, 5\} \leftrightarrow \{0, 6\}$ about its leftmost anchor. Therefore

$$F_{26}(X) = \{0, 6, 500, 506\}, \quad F_*(X) = \{1, 6, 501, 506\}. \quad (135.1)$$

This is a direct finite local-word calculation. It does not execute a CA simulator. The input has four particles and no valid marker triple; it is outside the admissible five-particle simulation. Equation (135.1) rules out identifying the two full-shift maps merely by changing a radius constant.

[write] Since $F_*(X) \neq F_{26}(X)$, the sparse evaluator of Report 27 and the quartic certificates of Report 28 (Part VI), which are for $F_{26} = F_{\text{new}}$, neither evaluate nor certify F_* (Section 137).

136 A separate complete radius two certificate

This section is independent of the two-scale construction. It concerns a uniform binary rule $f : \{0, 1\}^5 \rightarrow \{0, 1\}$ on the symmetric neighborhood $[-2, 2]$.

Theorem 136.1 (Certified radius two classification). *If the CA with local rule f conserves particle number on finite configurations and is injective on the full binary shift, then*

$$f(x_{-2}, x_{-1}, x_0, x_1, x_2) = x_j \quad \text{for some } j \in \{-2, -1, 0, 1, 2\}.$$

Thus only the five coordinate shifts satisfy both requirements at radius at most two.

The proof combines a complete finite parameterization of conservative tables with 423 explicit collision witnesses. The known count of 428 conservative five-input binary tables is recorded in Fuk s and Sullivan [34]; it is not a novelty claim. The fresh independent audit additionally derived a graph parameterization, reproduced all 428 tables and checked the supplied witnesses without running either original appendix program. We give that self-contained completeness argument here.

[write] The conservation criterion of Lemma 136.2 below is the radius-two analogue of the potential on the four-vertex de Bruijn graph with which SMC source 16 (Report 11, Part VI of `signal-machine-collision-certificates`) certifies the five conservative elementary rules (`$smc:su:sec:checks`); whether the lemma’s exact form is in the literature on number-conserving automata cited in Section 98 was not checked, and no attribution is made (Section 149, item 3). Theorem 136.1 is new in the collection; Report 70 claims no priority for it. The “fresh independent audit” is shipped as `70-two-scale-audit-r2-AUDIT.md`, the original appendix programs as `code/70-two-scale-pp-appendix-radius2_algebra.py` and `code/70-two-scale-pp-appendix-verify_radius2_algebra.py`.

136.1 Conservation is an integer graph potential

The binary order-four de Bruijn graph has sixteen four-bit vertices. A five-bit word $abcde$ is an edge from $abcd$ to $bcde$, with weight

$$g(abcde) = f(abcde) - c.$$

The graph is strongly connected: appending the four bits of a target vertex reaches it from any vertex.

Lemma 136.2 (Conservation criterion). *The binary table f is number-conserving on finite configurations if and only if there is an integer potential p on the vertices, normalized by $p(0000) = 0$, such that*

$$f(abcde) = c + p(bcde) - p(abcd) \tag{136.1}$$

on every edge.

Proof. Conservation first forces $f(00000) = 0$, since the zero configuration must remain zero. Fix a periodic binary word and place M copies between infinite zero tails. Except at a bounded number of boundary sites, the difference of output and input particle counts is M times the sum of g over one period. Boundary contributions are bounded independently of M . Conservation gives $0 = M\delta + O(1)$ for all positive M , so $\delta = 0$.

Every directed closed walk of the graph gives overlapping windows of a periodic binary word, so its total weight is zero. Define $p(v)$ as the weight of a directed path from 0000 to v . If two paths are chosen, append the same directed return path to 0000. Both closed walks have zero weight,

and the path weights coincide. Thus p is well-defined and integer-valued. Appending one edge proves (136.1).

Conversely, the zero self-loop in (136.1) gives $f(00000) = 0$. For a finite configuration, summing output minus input bits telescopes the differences of consecutive four-bit-window potentials. Both tails eventually equal 0000, so the sum is zero. This proves finite number conservation. $\square$

136.2 A complete finite enumeration

A spanning tree of the underlying undirected graph has fifteen edges. Choose a Boolean output of f on each tree edge. Its directed difference $f - c$ uniquely determines every potential from $p(0000) = 0$, using the negative difference when traversing the edge backward. Equation (136.1) then forces all thirty-two outputs. Retain a choice exactly when all forced values are Boolean.

There are exactly $2^{15} = 32\,768$ choices. Every conservative table occurs because its own tree outputs recover its potential. Different retained choices cannot give the same table because the chosen tree-edge outputs themselves differ. Every retained table is conservative by Lemma 136.2. Hence this enumeration is exhaustive without using a prior census as an assumption.

The independently written checker uses the following five-bit edge indices, in tree traversal order:

$$1, 16, 2, 3, 8, 24, 5, 18, 6, 7, 20, 28, 11, 13, 15.$$

Words are indexed in binary order 00000, ..., 11111. Six edges are traversed against their directed orientation. The recorded finite calculation rejects 32 340 non-Boolean assignments and retains 428 distinct complete tables, with no duplicates. The package includes the sorted table set and an integer potential for every retained table.

The original appendix reaches the same census from the leading-zero conservation identity

$$f(x_1, \dots, x_n) = x_1 + \sum_{k=1}^{n-1} [f(0^k, x_2, \dots, x_{n-k+1}) - f(0^k, x_1, \dots, x_{n-k})] \quad (136.2)$$

with $n = 5$ and $f(00000) = 0$. It enumerates the other fifteen leading-zero outputs. A separately written original verifier uses the spatially reversed identity. The graph audit above supplies a third, independently derived completeness route. Original programs and their receipts are retained as inert evidence; the Report 70 build and release tools execute none of them.

136.3 All nonprojections have explicit collisions

For each of the 423 nonprojection tables, the JSON certificate supplies a word length ℓ , distinct binary words a, b of that length, and a common output word y . For each cyclic index i , it records the equations

$$f(a_{i-2}, a_{i-1}, a_i, a_{i+1}, a_{i+2}) = y_i = f(b_{i-2}, b_{i-1}, b_i, b_{i+1}, b_{i+2}).$$

Repeating a and b bi-infinitely gives distinct full-shift inputs with equal images. Fundamental period need not equal the displayed word length. Each record is therefore a proof of noninjectivity, not a heuristic that a long trajectory appears irreversible.

The fresh graph audit checks table membership, all integer types and ranges, distinct inputs, uniqueness of the 423 rule records, every collision equation, and that the records together with the five projections cover the independently enumerated 428 tables. Its 4062 local bit equalities are exactly

$$2(164 \cdot 4 + 179 \cdot 5 + 80 \cdot 6).$$

The histogram of first witnessed lengths is 164 at length four, 179 at length five and 80 at length six. A separate bounded check in that audit exhausts all words of lengths one through six, confirming these are the first collision lengths for the corresponding rules. Only the listed collision equalities and complete table coverage are needed for Theorem 136.1.

The remaining tables are explicitly:

Coordinate offset j in $F(x)_i = x_{i+j}$	Rule integer
−2	4 294 901 760
−1	4 278 255 360
0	4 042 322 160
1	3 435 973 836
2	2 863 311 530

The output on input word index w is bit w of the rule integer. Each map has inverse coordinate offset $-j$; particle displacement has the opposite sign to j . All are full-shift bijections. This completes the finite-certificate proof.

136.4 What the classification does and does not imply

A nontrivial computationally universal binary, uniform, time-independent, number-conserving, full-shift-reversible CA in this symmetric-radius convention must therefore have radius at least three. This does not construct a radius-three universal example, classify larger neighborhoods or close the gap to the sufficient upper bound in (127.1). It does not apply unchanged to partitioned, time-dependent, second-order, nonuniform or larger-alphabet models.

Schranko and de Oliveira’s 2011 paper [35] studies reversible number-conserving rule derivation and representation. A bounded renewed search verified its bibliographic entry on the author’s publication list but did not obtain its full text; the publisher route was unavailable. Accordingly this report does not attribute a precise five-input classification to that paper or claim priority over it. The result here is an independently certified project appendix. Morita’s 96-state four-neighbor construction [4] is a distinct model and supplies no binary five-particle radius bound.

137 Dependencies certificate transport and attribution

137.1 What transports from earlier reports

The exact admissible agreement theorem transports semantic computation statements: an encoded source computation, its observer event, physical time or clean target has the same meaning under F_* . The transport must explicitly use admissibility of all involved states and the original certificate’s hypotheses.

It does not transport an arbitrary-input evaluator or verifier. Report 19’s evaluator implements an earlier ordered full-shift rule, and Report 26’s program implements its own parallel rule. Equation (135.1) already separates F_* from the latter. Old ordered-factor event budgets, malformed-state arithmetic transition fibers, Diophantine circuits and circuit-size counts cannot be relabelled as new-rule results. An arbitrary-input verifier for F_* would have to encode the complete raw recognition, same-anchor mutual uniqueness, isolation, hypothetical recognition test and simultaneous update. None is supplied here.

[write] Report 19’s evaluator is Part IV; Report 26’s program is `code/26-parallel-involutions-parallel_particles.py`. The same holds for Report 27’s sparse evaluator and Report 28’s quartic certificates, which are for Report 26’s rule.

The construction is a finite mathematical definition of a local Boolean rule. This gives effective existence in principle, but not an implemented new-rule evaluator, practical universal evaluation, small arithmetic circuit or template-independent running-time bound. The sufficient radius does not measure the work needed to evaluate its enormous local truth table.

137.2 Established architectural precedents

Marker automorphisms, conserved recognition data and products of involutions predate this construction. The bounded primary-source audit retained from Report 26 supplies the following precise attributions.

Boyle, Lind and Rudolph’s marker replacements implement simultaneous local data permutations; their Theorem 2.6 constructs a product of two marker involutions [22]. Arbitrary data permutations need not conserve particle number. Equal-weight permutations provide a conservative specialization, which is an observation rather than a separate theorem attributed to that paper.

Toffoli and Margolus discuss conserved-landscape permutations in Section 5.3 and conservative local block operations in Section 5.5 [23]. The landscape must remain identifiable to undo the update. Their illustrative binary flip is not itself number-conserving. Section citations refer to the inspected corrected author-uploaded reprint in the inherited audit, not newly verified original pagination.

Maldonado, Moreira and Gajardo’s Proposition 2 uses invariance of recognized left, right and malformed classes to extend a reversible partitioned CA [25]. That recognition-preservation proof is particularly close conceptually. Their radius-one pair-update map fixes its malformed class; this assertion does not describe the entire composition with its shift factor. Their universal time-symmetric conservative constructions use a larger alphabet, and their time-reversing involution need not itself conserve numeric mass.

Gajardo, Kari and Moreira’s Proposition 1 identifies time-symmetric CAs with products of two CA involutions [24]. The identities in Section 133 instantiate that standard algebra. A reversible CA is not automatically a product of two individually conservative binary CA involutions; here those properties are proved for each factor.

The specific quantitative statement proved in this report separates recognition radius a from control radius r , obtains (128.5), and shows that its guard-free anchor predicate preserves every intended edge of this literal compiler. Neither the architectural ideas nor the lower-radius appendix are presented as externally new. The bounded search and inaccessible 2011 full text do not support a priority claim.

138 Evidence and reproducible presentation

The companion release retains the complete frozen proof packet, its original independent proof audit, the complete fresh static and radius-two audit directories, and the original Report 26 archive as an *inert* ZIP. The original archive contains scientific code and historical execution receipts. It is preserved for authentication and provenance, not invoked by the report tools.

The frozen proof packet has external manifest SHA-256

```
f1c19ca86641cbf5ebaed81d0c78e82563258753bbf8f38f2f46229247b71fa0
```

and its authenticated Report 26 archive has SHA-256

```
20a23b1ee22aed461942da4def6bade2bd55d777fce1fe49b6ff83248b7442e4.
```

The fresh combined audit manifest is pinned by

`3bfe65db1ad9c0aab9fd13641cfbad843a292d2c2dbec444fa0f7ebb856876b1`.

Its accepted proof audit independently checks the full-shift lemma, literal templates, every signed boundary, resource arithmetic, malformed witness and inherited limits. Its radius-two audit derives the graph criterion, enumerates all tables and directly checks every collision. The source and audit files are copied without changing bytes, modes or modification times; their full copied inventories are recorded separately.

The mathematical proof is not proof-assistant formalized. During Report 70 authoring, no archived or upstream scientific program, source-machine interpreter, trajectory simulator, saved schedule or Lean program is executed. The only report executions are presentation and release operations. The fresh audit's independently authored static checks are already completed evidence, not newly rerun author programs.

The editable flattened `Report70.tex` matches the modular manuscript. The Report70-owned tools adapt the fully inspected accepted Report 68 presentation contract, with the corresponding Report 66 sources also compared. A build uses an allowlisted environment, a fresh local TeX format and font map, disabled shell escape, three passes, and the union of the format and every pass's recorded inputs. A locked replay checks interpreter and executable hashes, all recorded system TeX/font inputs and selected maps. Its lock does not inventory shared libraries, Python's standard library or the operating system. Reproducible typesetting is evidence about the presentation, not a theorem prover.

[write] Report 68 (*Gaussian fluctuations of encoded gap counts and inverse ranks*) was never delivered to ProveIt: only its README survives, as `38-low-arity-qa-pred-tools-README68.md` in `signal-machine-collision-certificates`, delivered inside the archive of Report 69. Report 66 is placed in the same report (prefix `36-bounded-cert-`). Report 70's copies of their READMEs and presentation tools are not shipped here. What can be checked of the sentence above is recorded in Section 149, item 1. Of the companion release, this report ships the proof packet, the two fresh audits, the input pins and the review and build records; the Report 26 archive and its readable L^AT_EX and PDF are the committed archive of `db37d18c8` and its members, and every member of the frozen-packet ZIP is shipped (README).

The release manifest records the exact descendant inventory, bytes, modes and nanosecond modification times. Its digest must be obtained outside the archive; an internally self-consistent manifest is not an external signature. The tools reject extra entries, symlinks, special files, hardlinks and unsafe output aliases. The supplied extractor restores authenticated descendant modification times that ordinary ZIP extractors may discard. Manifest-file metadata and the newly created top-level root's metadata are outside the inventory boundary. Read-access times and original creation/change times are not transported as integrity claims.

The README gives authentication, locked-build and deterministic-archive commands with the final manuscript and dependency pins. The original science and independent audit programs remain inert during all those operations. Automated PNG structure checks are distinct from the required visual inspection of every rendered page. Independent manuscript and release-tool review records are retained with the completed package.

139 Further mathematical directions

The remaining large scale comes primarily from unary gap encoding of the signed head modes: $D = 2m + 4p$. Reducing the source mode count or replacing that finite-mass head encoding

would be a different construction and could change exact clocks. The present result improves the coefficient without changing those objects.

Two immediate technical directions remain. First, type-dependent geometric separation might reduce a common exclusion bound further, but any argument must protect all isolated rejected controls and exclude cooperative recognition births, not merely keep current writes disjoint. Second, a practical evaluator or arithmetic certificate should prove complete candidate discovery on arbitrary malformed multihead inputs; valid-orbit simulation tests alone would not establish that contract. Formalizing the anchor and signed-corridor arguments could help prevent confusion between current and predecessor head anchors or between current and old marker anchors.

The completed conclusion is narrower and exact: the same five-particle admissible dynamics admits a new conservative binary full-shift extension, composed of two involutions, with the sufficient radius in (127.1). The universal number remains an inherited-ledger substitution and the radius-two result remains an independent finite classification.

Shorter exactness windows for the binary five particle compiler

A companion to Report 70

Research Report 71

4 October 2026

Abstract of Report 71, as delivered. The conservative binary five-particle compiler admits a further reduction in its sufficient radius. Keep pair exactness $L = 3D + 4$, but shorten triple exactness from $3b + 1$ to $b = 4D + 5$ in both its edge and phase blocks. We define both new blocks on every bi-infinite binary input, prove that each is a number-conserving involution, and prove all-anchor recognition equivalence and exact forward and inverse agreement on the entire admissible doubled micrograph. The resulting rule has sufficient radius $76D + 105 + 3J$, where $D = 2m + 4p$ and J is the source's finite counter-class cut. A changed-output-support argument improves this to $76D + 104 + 3J$ for the very same rule; the boundary-read predicates are unchanged.

The inherited universal-source ledger gives 38 722 713, or 38 722 712 with the support refinement. Its source table remains absent and universality remains an inherited premise. Explicit malformed inputs distinguish each new block and the composition from both predecessors. Pair-padding counterexamples concern only the unchanged templates and selector with an enabled nonzero branch. There is no new-rule evaluator, arbitrary-input arithmetic certificate, minimum-radius or priority claim.

[write] Delivered as `Shorter_exactness_windows_sources.zip` (batch 91, manuscript 23; arrival 0d7f51c44), dated 4 October 2026; files placed by 22a8ca89e with the prefix `71-short-w` `indows-`. No author is named and no ProveIt commit is pinned. Report 71 pins Report 70's release manifest (0f96f8aa8bba...) and PDF (d061fda50f26...); both are the files of the delivered Report 70 archive (checked at placement and again for this write), so the predecessor it pins is the text printed just before it. It also pins Report 26's archive, as Report 70 does. Section 141 restates Report 70's source interface, templates and admissible graph (Sections 130–131) with the new exactness predicate (141.12); the safety and agreement proofs of Sections 142–143 follow Report 70's pattern for the shorter windows, and Report 70's radius-two classification is not used. Report 71's $\backslash N$ is $\backslash N_{ge}$ here; its macros $\backslash Es$, $\backslash Ps$, $\backslash Fs$ (the subscript s stands for “short”, not for the side s of the phase triples) and $\backslash Cand$ print as delivered. The text printed is the delivered third version; the second differed only by the words “For a halting run” in the definition of Θ (Section 144; shipped `data/71-short-windows-qa-V3-HALTING_QUALIFIER.diff`).

140 Result and comparison

Write $I_t(u) = [u - t, u + t] \cap \mathbb{Z}$, with both endpoints included. A configuration is an arbitrary $x \in \{0, 1\}^{\mathbb{Z}}$. Number conservation means preservation of the ordinary number of ones on finite-support inputs; no divergent infinite sums are compared. Report 26 and Report 70 give full-shift extensions of a common five-particle simulation. This report changes *both* blocks of that extension. It does not correct either predecessor's definition or claimed sufficient bound.

Theorem 140.1. *For every finite source satisfying Section 141, the maps defined below give a binary CA $F_s = P_s \circ E_s$. Both factors are full-shift, number-conserving involutions. On every natural-counter home, both signs of every strict intermediate of an enabled source-edge subdivision, and every translate, they equal the common predecessor edge matching and sign flip. Consequently both forward and inverse dynamics and exact clocks are unchanged. Moreover,*

$$\begin{aligned} R(E_s) &\leq 13b + 10 + 3J, & R(P_s) &\leq 6b, \\ R(F_s) &\leq 19b + 10 + 3J = 76D + 105 + 3J. \end{aligned} \tag{140.1}$$

For the same rule, $R(P_s) \leq 6b - 1$ and $R(F_s) \leq 76D + 104 + 3J$.

Construction	Triple exactness	Total sufficient bound	Inherited ledger
Report 26	$3b + 1$	$180D + 258 + 9J$	91 711 698
Report 70	$3b + 1$	$108D + 149 + 3J$	55 027 013
Report 71, main	b	$76D + 105 + 3J$	38 722 713
Same rule, support refinement	b	$76D + 104 + 3J$	38 722 712

Table 16: Bounds, not measured or minimum radii. Pair exactness remains L in every row. The final column substitutes $D = 509\,508$, $J = 0$ and inherits the same universal-source premise.

Report 70 supplies broader architectural background, attribution, clean-target details and a separate radius-two classification. That classification is not repeated or used here. The complete definitions and proof obligations for the present theorem follow.

[write] Report 26's bound is Theorem 89.1, Report 70's Theorem 127.1. For the smallest source of Report 26's comparison (two controls and one zero-update branch, $D = 4$, $J = 0$; Section 95) the four rows of the table give 978, 581, 409 and 408, against the ordered rule's 540.

141 Source, geometry and complete endpoint families

141.1 Source interface and admissibility

The source is a finite deterministic partial injection on $Q \times \mathbb{Z}_{\geq 0}^2$. A branch $e : q \rightarrow q'$ selects side $v \in \{-1, +1\}$ and its corresponding counter index j (left for $v = -1$, right for $v = +1$), updates that counter by $\Delta \in \{-1, 0, +1\}$, and has domain predicate G_e . Enabled outputs are natural. Branch domains at each control and actual branch images at each target control are disjoint on *all* natural-counter IDs, including unreachable ones. The exact image predicate is

$$I_e(c') = G_e(c' - \Delta e_j) \wedge (c' - \Delta e_j \in \mathbb{Z}_{\geq 0}^2).$$

Both G_e and I_e depend only on classes $0, 1, \dots, J, > J$ of the two counters. There is a designated halt control with no outgoing branch. Other controls may be stuck; such blocking is not designated halting.

Let $m \geq 1$ count controls, p nonzero-update branches, and α zero-update branches. The unsigned modes are homes H_q and outbound/inbound modes O_e, I_e for each nonzero branch. Give all signed modes distinct gaps in $\{1, \dots, D\}$, and set

$$\begin{aligned} D &= 2m + 4p, & S &= 2D + 2, & B_2 &= D + 1, & L &= 3D + 4, \\ b &= L + D + 1 = 4D + 5, & Z &= 10b + 10 + 2J, & r &= Z + J. \end{aligned} \tag{141.1}$$

Thus $D \geq 2$, $J \geq 0$ and $r \geq 3b$. A head in signed mode u^σ at x is the pair $\{x, x + d(u^\sigma)\}$. Its anchor is always its leftmost particle, including during leftward motion. At origin zero a home is

$$\text{Enc}(q, c_\ell, c_r)^\sigma = \{-(Z + c_\ell), 0, Z + c_r, S, S + d(H_q^\sigma)\}.$$

The three markers have separations at least Z . The encoding uses exactly five particles over a binary zero background.

An enabled zero-update edge is one unsigned home-to-home microedge. For a nonzero edge, let o be the origin, m_0 the selected old marker, and $m_1 = m_0 + v\Delta$. Dispatch changes the home head at $o + S$ to O_e at $o + vS$. Its anchor travels by v to $m_0 - vS$. The endpoint interaction moves the

marker to m_1 and changes the head to I_e at $m_1 - vS$. The head travels by $-v$ to $o + vS$, then commits to $H_{q'}$ at $o + S$. Dispatch reads G_e on old counters; commit reads I_e on post counters. The admissible unsigned graph consists of every natural-counter home and *only* strict intermediates of enabled subdivisions. Guard-invalid moving geometry is excluded. Branch-specific modes distinguish intermediates, and exact image disjointness prevents mergers at homes, so this graph is a partial injection. Double it by signs: each unsigned edge $c \rightarrow c'$ gives an edge $c^+ \leftrightarrow c'^-$; phase pairs $c^+ \leftrightarrow c^-$. Include all translates.

141.2 Literal edge endpoints

All displayed sets are offsets from the invariant template anchor zero. For a moving mode u , write $d_+ = d(u^+)$, $d_- = d(u^-)$, and use direction $w = v$ for O_e , $w = -v$ for I_e . The free pair and near-travel triples are

$$\{0, d_+\} \longleftrightarrow \{w, w + d_-\}, \quad (141.2)$$

$$\{0, wt, wt + d_+\} \longleftrightarrow \{0, w(t+1), w(t+1) + d_-\}, \quad S \leq t \leq L, \quad (141.3)$$

$$\{0, -wt, -wt + d_+\} \longleftrightarrow \{0, -w(t-1), -w(t-1) + d_-\}, \quad S+1 \leq t \leq L. \quad (141.4)$$

For each nonzero branch, the interaction templates are

$$\{0, S, S + d(H_q^+)\} \longleftrightarrow \{0, vS, vS + d(O_e^-)\}, \quad (141.5)$$

$$\{0, -vS, -vS + d(O_e^+)\} \longleftrightarrow \{v\Delta, v\Delta - vS, v\Delta - vS + d(I_e^-)\}, \quad (141.6)$$

$$\{0, vS, vS + d(I_e^+)\} \longleftrightarrow \{0, S, S + d(H_q^-)\}. \quad (141.7)$$

Their guards are respectively G_e , none, and I_e . For a zero-update branch use

$$\{0, S, S + d(H_q^+)\} \longleftrightarrow \{0, S, S + d(H_q^-)\}, \quad G_e = I_e. \quad (141.8)$$

Each branch type remains distinct, even when several branches share a geometric home endpoint. There is no priority ordering.

141.3 Literal phase endpoints and exactness

The phase family contains, for each moving mode, the pair and both sides of the near triples

$$\{0, d_+\} \longleftrightarrow \{0, d_-\}, \quad (141.9)$$

$$\{0, st, st + d_+\} \longleftrightarrow \{0, st, st + d_-\}, \quad s \in \{-1, +1\}, \quad S \leq t \leq L, \quad (141.10)$$

and for each control the home triple

$$\{0, S, S + d(H_q^+)\} \longleftrightarrow \{0, S, S + d(H_q^-)\}. \quad (141.11)$$

There is no free-home phase type and no contextual phase guard.

For either family let P_g, Q_g denote the displayed endpoints and put $B_g = B_2$ for a pair, $B_g = b$ for a triple. The local map $\tau_{g,u}$ exchanges their complete words on $I_{B_g}(u)$, fixes every other word there, and fixes all sites outside that interval. The supports fit: free edges lie in $[-1, D+1]$; behind edges in $[-L-1, L+D+1]$; ahead edges in $[-L, L+D]$; home interactions in $[-S, S+D]$; endpoint interactions in $[-S-1, S+D+1]$. Each lies in its declared window. Phase supports satisfy the sharper bounds in Section 144.1. Distinct signed gaps make the two words distinct; their weights are equal. Thus every $\tau_{g,u}$ is an everywhere-defined conservative involution.

The *new* geometric predicate in both families is

$$\ell_g = \begin{cases} L & \text{pair,} \\ b & \text{triple,} \end{cases} \quad \rho_g(x, u) \iff \{i \in I_{\ell_g}(u) : x_i = 1\} \in \{u + P_g, u + Q_g\}. \quad (141.12)$$

In particular, the pair writes only $I_{B_2}(u)$, not the common $I_b(u)$. Arbitrary bits outside its small write window are preserved, including particles at distances between L and b . Zero-filling the common window would define a different, incorrect rule.

For edge types let c_g be ρ_g conjoined with its contextual guard, when present. For each counter, the guard reads $u \pm (Z + k)$, $0 \leq k \leq J$: exactly one occupied read selects that finite class, no occupied read selects $> J$, and multiple occupied reads reject the guard. These conventions are finite Boolean definitions on malformed inputs too. All class sites lie outside every rewrite. Hence c_g is invariant under its own $\tau_{g,u}$, and it reads radius at most r .




Figure 9: Schematic window comparison, not a full configuration or a scale drawing. On admissible matches, remote markers lie beyond both windows. On malformed inputs an extra particle between the two boundaries can change recognition; Section 145 gives exact examples.

142 Full-shift definition and safety of both new blocks

142.1 Edge: mutual guarded uniqueness and raw-anchor selection

At every integer anchor use the *full* guarded type set

$$C_u(x) = \{g \text{ in the edge family} : c_g(x, u)\}.$$

Define

$$T_u(x) = \begin{cases} \tau_{g,u}x, & C_u(x) = \{g\} = C_u(\tau_{g,u}x), \\ x, & \text{otherwise.} \end{cases} \quad (142.1)$$

Both uniqueness tests include all types. If the first case acts, the same unique type passes in the image and reverses it; otherwise the input is fixed. Thus $T_u^2 = \text{id}$ even with zero, multiple, conflicting or newly enabled types. It writes only $I_b(u)$ and preserves weight there. Its replacement word is determined by $I_r(u)$: hypothetical tests use already known replacement bits and otherwise unchanged input bits in that same window, adding no radius.

Coalesce geometric types only at the raw-anchor level:

$$\mathcal{R}(x) = \{u \in \mathbb{Z} : \rho_g(x, u) \text{ for some edge type } g\}.$$

Recognition reads radius $a = b$, since $L \leq b$. Nonidentity T_u implies $u \in \mathcal{R}(x)$. Set

$$H = \max\{2(b + a), b + r\} = b + r. \quad (142.2)$$

Select u if $u \in \mathcal{R}(x)$, no other $v \in \mathcal{R}(x)$ has $|v - u| \leq H$, and

$$\mathcal{R}(T_u x) = \mathcal{R}(x). \quad (142.3)$$

Apply all selected T_u simultaneously to define E_s . The equality compares *raw integer anchors*, not guarded types or already eligible anchors. Only anchors within $b + a = 2b$ of u can change recognition. Their recognition words lie in $I_{b+2a}(u) = I_{3b}(u)$. Together with T_u 's control word this makes the prospective test a finite radius- r test.

Lemma 142.1. *The map E_s is a number-conserving involution on $\{0, 1\}^{\mathbb{Z}}$, with radius at most $3b + r = 13b + 10 + 3J$.*

Proof. Distinct selected centers are farther apart than H and have disjoint write intervals. If a radius- b recognition window met writes at two selected centers u, w , then $|u - w| \leq 4b \leq H$, a contradiction. Every post-update recognition word therefore equals its old word or the word after exactly one selected T_u . That update passed (142.3); hence

$$\mathcal{R}(E_s x) = \mathcal{R}(x)$$

at every integer anchor on every bi-infinite input. This covers initially absent anchors, dense malformed configurations, and potential cooperative births or deaths.

Now take any recognized isolated anchor u , whether selected or rejected. Every other selected w has $|w - u| > H = b + r$, so its write misses $I_r(u)$, which contains all control and prospective data. The complete status analysis is:

1. An isolated rejected anchor's determining word is unchanged, so its failed prospective test remains failed.
2. A selected anchor sees precisely its own T_u image. The equality $T_u^2 = \text{id}$ interchanges the two recognized sets in (142.3), preserving success.
3. A nonisolated recognized anchor stays nonisolated by recognition-set invariance; it remains blocked even if its guarded types change.
4. An absent anchor remains absent by the same invariance.
5. A recognized isolated identity control passes prospectivity harmlessly. Its protected radius- r word prevents a new control from appearing there on the second application.

Thus the *entire* selected-anchor set is invariant. Each selected control sees only its own prior image on reapplication and undoes its disjoint change. Therefore $E_s^2 = \text{id}$.

Equal-weight changes on disjoint blocks conserve finite particle number. A finite-support input permits only finitely many nontrivial changes, since each uses an original occupied site in its disjoint block; finite support is preserved. No total on infinite inputs is needed.

For locality, a changed output coordinate has a writing anchor within b . Isolation reads $H + b$ about that anchor and dominates radius- r control/prospectivity. Thus $R(E_s) \leq b + (H + b) = 3b + r$. Finite Boolean tests and translation covariance make this a binary CA. Closed intervals and exclusion at distances $\leq H$ are used throughout. $\square$

142.2 Phase: typed keys and its own prospective filter

This is a new phase block, not the former phase map with a relabelled radius. Keep the type-anchor keys

$$\mathcal{C}_P(x) = \{(g, u) : g \text{ in the phase family and } \rho_g(x, u)\}.$$

Endpoint orientation is not part of a key, but type is. Put $r_P = b$ and $H_P = 2(b + r_P) = 4b$. Select $k = (g, u)$ if it is a candidate, no *other* candidate key has anchor within H_P of u (including another type at u), and

$$\mathcal{C}_P(\tau_k x) = \mathcal{C}_P(x). \tag{142.4}$$

Apply the selected transpositions simultaneously to define P_s . The equality requires only all phase types at anchors within $b + r_P = 2b$ of u , so it reads $I_{b+2r_P}(u) = I_{3b}(u)$.

Lemma 142.2. *The map P_s is a number-conserving full-shift involution with radius at most $6b$.*

Proof. Selected writes are disjoint. A radius- b candidate window cannot meet two of them, because then their centers would be at distance at most $4b$. Each post-update candidate predicate consequently sees no rewrite or exactly one individually prospectively stable rewrite. Thus the *complete typed candidate set* is invariant, including initially absent keys and keys of other types.

For every isolated candidate, any other selected write misses its radius- $3b$ prospective-decision window: the center distance is greater than $4b$, the write radius is b . An isolated rejected key therefore keeps its failed prospective test. A selected key's own involution interchanges the sets in (142.4), preserving success. Nonisolated keys stay blocked and noncandidates stay absent. Hence the entire selected-key set is invariant and a second application reverses the same disjoint involutions. Equal weights give finite conservation and finite support as in Lemma 142.1.

Isolation reads $H_P + r_P = 5b$ about a candidate anchor; prospectivity reads $3b$. A possibly writing anchor lies within b of an output site, so $R(P_s) \leq 6b$. No uniqueness theorem is assumed on malformed phase inputs: multiple same-anchor types are competing keys and are rejected by typed isolation. This differs deliberately from raw-anchor coalescing in the edge block. $\square$

143 All-anchor admissible agreement

Full-shift safety alone does not establish simulation preservation. We now show that shortening exactness creates no admissible match at *any* integer anchor, then classify every intended edge and phase case.

143.1 The close pair and the only possible singleton

Every admissible configuration has exactly one occupied pair at distance at most D , the head. Head-marker distances are at least $S - D = D + 2 > D$, and marker-marker distances are at least Z . Every literal endpoint likewise contains a unique close pair; every triple's singleton has distance greater than D from both head sites. This includes the reverse behind endpoint at distance $L + 1$ and the shifted endpoint interaction.

Thus any new raw match, at any anchor, must use the actual head pair. Its distinct signed gap fixes the signed mode and endpoint orientation; any displayed singleton must be an actual marker. Every triple singleton lies within $L + 1$ of the actual head anchor: near-travel distances lie in $S, \dots, L + 1$, and all interaction distances are S . Since

$$Z > 2(L + 1),$$

at most one actual marker can support *any* triple using this head. This common head-to-marker bound excludes cross-template competition; bounding only one template's diameter would not suffice.

For a fixed endpoint and actual head pair, the offsets force the invariant anchor. Ordinary triples anchor at the singleton marker. Only the reverse endpoint uses an anchor one site from its current singleton, as detailed below. Other markers are therefore at least Z from an ordinary anchor, or $Z - 1$ from the old endpoint anchor. Since

$$Z - 1 > 3b + 1 > b,$$

they lie outside both old and new triple exactness windows. All three displayed particles fit in I_b . Consequently a geometric triple match on an admissible state satisfies old exactness $3b + 1$ if and only if it satisfies new exactness b . Pair predicates are literally unchanged. This proves all-anchor, all-type old/new geometric recognition equivalence; the following classification makes the boundaries and guards explicit.

143.2 Homes and multiple raw branch types

At H_q^+ only outgoing dispatch/direct forward endpoints contain its gap, all at the origin. At H_q^- only incoming commit/direct reverse endpoints contain its gap, again at the origin. No moving free or travel template uses a home gap. Hence raw edge anchors are empty or a singleton, although several raw types may occur there.

At plus homes the full guarded set is empty or a singleton by source-domain disjointness; at minus homes it is empty or a singleton by *exact actual-image* disjointness. Class reads are faithful: counter markers encode the classes while the head lies inside $b < Z$. The reverse image guard includes existence of natural pre-counters and therefore enters only an enabled subdivision. Replacing it with the forward guard would not justify this conclusion. If there is no enabled incident edge, every control T_u is identity even when raw branch types occur at the origin. Direct zero-update edges retain $G_e = I_e$.

143.3 Both signed corridors and every equality case

Let w be a moving mode's direction, A its departure marker and C its destination marker. With head anchor x , write

$$N = w(C - A) \geq Z, \quad t = w(x - A), \quad S \leq t \leq N - S.$$

Outbound has $w = v$ and $N = Z + c$; inbound has $w = -v$ and $N = Z + c + \Delta \geq Z$, since enabled post-counters are natural. The third marker is remote. Since $N > 2(L + 1)$, the intervals below are disjoint and exhaustive.

For a plus head the free key is x , so pair exactness requires $L < t < N - L$. For a minus head the free key is $x - w$, *not* x , so pair exactness requires $L < t - 1 < N - L$. Converting to integer ranges gives:

Sign	Incidence	Range of t	Invariant key
plus	behind travel	$S \leq t \leq L$	A
plus	free travel	$L + 1 \leq t \leq N - L - 1$	x
plus	ahead travel	$N - L \leq t \leq N - S - 1$	C
plus	forward interaction	$t = N - S$	as below
minus	reverse interaction	$t = S$	as below
minus	behind inverse	$S + 1 \leq t \leq L + 1$	A
minus	free inverse	$L + 2 \leq t \leq N - L$	$x - w$
minus	ahead inverse	$N - L + 1 \leq t \leq N - S$	C

Table 17: Complete raw edge partition. The head remains leftmost within its pair for $w = -1$ too; the two particles are not reflected in order.

In particular, a behind triple takes plus departure distance L to minus distance $L + 1$. At the image, the free predecessor key still sees A at distance L , so inverse free recognition fails. At the ahead boundary, a free step takes plus arrival distance $L + 1$ to minus distance L . The

same predecessor key sees C at distance $L + 1$, so inverse free recognition succeeds; the inverse ahead-triple range ends at current arrival distance $L - 1$ and cannot compete.

At plus arrival distance S , travel has already ended and only the forward interaction remains. At minus departure distance S , inverse travel has not begun and only the reverse interaction remains. The adjacent $S + 1$ and $N - S - 1$ positions are exactly the listed travel cases. These are signed algebraic statements valid for $w = \pm 1$. Within any triple range, the signed gap, unique possible marker, side and integer distance determine exactly one type. No extra match arises from shortening, by the all-anchor argument above.

143.4 Every interaction and the old endpoint key

There are four moving interaction incidences:

1. Outbound minus at departure is reverse dispatch, at the origin, with its true domain guard because the corridor is enabled.
2. Outbound plus at arrival is forward endpoint, at the old selected marker m_0 .
3. Inbound minus at departure is reverse endpoint, at the *old* key $m_0 = m_1 - v\Delta$, not the current marker m_1 .
4. Inbound plus at arrival is forward commit, at the origin, with its true exact image guard because the post-counters came from the enabled branch.

In the reverse endpoint word (141.6), the singleton offset is $v\Delta$ and the head-anchor offset is $v\Delta - vS$. Both independently force $u = m_1 - v\Delta = m_0$. The actual marker-head distance is still S ; only the invariant key is shifted. Other markers remain beyond $Z - 1 > 3b + 1$, so both exactness tests hold. This proves the old-key assertion for both physical sides $v = \pm 1$ and both update signs $\Delta = \pm 1$.

Branch-specific outbound/inbound gaps exclude all other branches' interactions. Within the same branch an actual singleton can fit the interaction only at its stated wall, where the corridor partition excludes travel. The third marker is outside the possible $L + 1$ singleton range. Hence every moving admissible state has exactly one raw type-anchor, with the required guard true.

143.5 The desired edge and phase selections all pass

For every intended doubled edge, its paired endpoint is again admissible. Both endpoints have the same guarded singleton $\{g\}$ at the same invariant key: predecessor head anchor for free travel, contacted marker for near travel, origin for dispatch/commit/direct, and old marker for endpoint interaction. Their entire raw-anchor sets are the same singleton. Therefore T_u is the intended swap and

$$\mathcal{R}(x) = \{u\} = \mathcal{R}(T_u x).$$

Mutual uniqueness, isolation and full prospectivity pass. At missing edges all controls are identities, so E_s fixes the state. This covers all source branches and counter sizes, both signs, every translate and both reflection ends.

At a home the sole phase candidate is its home type at the origin. At a moving state the free phase key is the unchanged head anchor x for *both* signs. It is a candidate exactly when every marker is farther than L from x . Otherwise a unique marker has distance in $S, \dots, L$, and its side and distance specify the sole near-phase triple. The close-pair and cross-template marker argument excludes every other candidate at every anchor; shortening introduces none.

A phase swap fixes every marker, the head anchor and the unsigned mode, changing only the signed gap. Its free/near/home case, type and key therefore do not change. Its full typed candidate set is the same singleton before and after, so typed isolation and prospectivity pass. Thus P_s is the exact sign flip on the entire admissible graph. This completes admissible agreement, including inverse agreement because the inverse composition is $E_s P_s$.

144 Radius, unchanged dynamics and exact scope

Lemmas 142.1 and 142.2 give

$$R(F_s) \leq R(E_s) + R(P_s) \leq 19b + 10 + 3J = 76D + 105 + 3J.$$

Both factors preserve the admissible set and equal the predecessors' restricted blocks, so every forward step and every inverse step is unchanged. Also

$$F_s^{-1} = E_s P_s, \quad E_s F_s E_s = F_s^{-1}, \quad P_s F_s P_s = F_s^{-1}.$$

These are two individually conservative involutions and time reversers.

144.1 A support refinement for the very same rule

Every phase endpoint lies in $I_{L+D} = I_{b-1}$. For a positive near endpoint the extremum is $t + d = L + D = b - 1$, including the equality $t = L, d = D$. The two negative-side head sites lie in $[-L, -S + D]$, and their singleton is at zero, so the full negative near support lies in $[-L, 0]$. Home endpoints lie in $[0, S + D]$ and free endpoints in $[0, D]$. Thus every complete phase endpoint word has zeros at both $-b$ and b .

On any non-endpoint word the transposition is identity; on either endpoint word the two boundary zeros remain zero. Therefore no phase transposition can change an output site at either boundary, on *any* input. The triple predicate nevertheless still reads those two zeros; the transposition still reads the complete I_b word. No predicate, transposition, isolation threshold $H_P = 4b$ or prospective test has changed.

Eligibility still reads radius $5b$ about its anchor. Only the anchors that can *actually change* a specified output site now lie within $b - 1$ rather than b . The output is determined by its old bit and those possible changes, proving for the same map

$$R(P_s) \leq 6b - 1, \quad R(F_s) \leq 19b + 9 + 3J = 76D + 104 + 3J. \quad (144.1)$$

This is an upper-bound refinement, not a smaller-window definition or a minimum-radius theorem. The simple common-support observation does not apply to edge endpoints: a positive reverse-behind endpoint at $t = L, d_- = D$ reaches $L + 1 + D = b$.

D	J	Edge bound	Phase, main	Total, main	Same-rule refinement
8	0	491	222	713	712
18	1	1 014	462	1 476	1 475
509 508	0	26 494 491	12 228 222	38 722 713	38 722 712

Table 18: Exact substitutions into sufficient bounds. The last row is conditional on the inherited universal-source ledger.

144.2 Clocks, observers, targets and dependencies

For a nonzero source edge at old selected counter c , the travel lengths remain $Z + c - 2S$ and $Z + c + \Delta - 2S$. Its three interactions give the unchanged physical clock

$$\tau_e(c) = 3 + 2(Z + c) + \Delta - 4S.$$

A zero-update edge still takes one step. The binary alphabet, five particles and anchored halt-observer length $3D + 3$ are unchanged. A missing graph edge still causes a one-step sign reflection.

For a halting run, let Θ be the original forward computation's total physical microedge time, the sum of its inherited clocks using the enlarged wrapper source's constants. The clean-target first time $2\Theta + 2$ and reflected period $4\Theta + 6$ transport with their original hypotheses. In particular, the clean-target bridge requires a *control-wide* predecessor-free initial normalization, not just one predecessor-free input ID, and all constants must be recomputed for that enlarged source. Its target restores the input counters and is input-dependent. Equivalence of positive return or periodicity with designated halting additionally requires no nonfinal blocking on the claimed source-input family; a nonfinal stuck home otherwise reflects and cycles too. Report 70 gives the full wrapper construction. [write] Report 70's wrapper is in Section 133; it is Report 15's (Section 13).

The endpoint-template count remains $8pD + 29p + m + \alpha$. Two mathematical involutions do not imply two cheap arithmetic factors or template-independent Boolean evaluation cost. Old source programs, malformed-input transition fibers, circuit counts, ordered-factor budgets, schedules, and arithmetic or Diophantine certificates do not become implementations or arbitrary-input certificates of this new full-shift map. Semantic admissible witnesses transfer through the proved agreement; malformed certificates do not transfer automatically.

The inherited universal receipt records $m = 122\,622$, $p = 66\,066$, $\alpha = 75\,495$, $J = 0$, hence $D = 509\,508$. Its source-table SHA-256 is

fa61d06178d718511232a80a05e02192d940fde3c273f0c202f0623253b2ade3.

That table is absent from the preserved Report 26 archive and this evidence packet. Its ledger, historical universality and required normalizations remain inherited premises. The source-uniform theorem does not require a universal table; the numerical substitution does not prove universality anew. No universal source, template array, new-rule evaluator or truth table was materialized or executed here.

145 Actual malformed inequalities for both blocks

Take ordered controls (q, h) , one true-guard right increment $q \rightarrow h$, and $J = 0$. Its exact image is right-counter positivity, expressible by classes $0, > 0$. Use the inherited literal gap order

$$H_q^\pm : 1, 2; \quad H_h^\pm : 3, 4; \quad O^\pm : 5, 6; \quad I^\pm : 7, 8.$$

Thus $D = 8$, $S = 18$, $B_2 = 9$, $L = 28$, $b = 37$, old triple exactness $K = 112$, and $Z = r = 380$. Write E_{26}, E_{70} for the predecessor edge blocks and $P_{26} = P_{70}$ for their shared phase block.

145.1 Triple shortening changes each block separately

Let $X = \{0, 18, 23, 80\}$. Its sole close pair is $\{18, 23\}$, gap five, so every match must use this O^+ head. Only singleton 0 is within $L + 1 = 29$ of head anchor 18; particle 80 is 62 away. The

free pair fails exactness because its radius-28 window contains 0. The only geometric edge type is behind O at $t = 18$, key 0; the only geometric phase type is near O on the positive side at distance 18, key 0. Both old triple predicates see particle 80 in $I_{112}(0)$ and fail. Thus all predecessor blocks fix X .

New triple exactness sees precisely $\{0, 18, 23\}$ in $I_{37}(0)$. The edge endpoint is $\{0, 19, 25, 80\}$, with O^- gap six. Its sole inverse-behind type has key 0 at current distance 19; the free inverse key 18 still sees marker 0. Particle 80 cannot support another triple. Both guarded uniqueness tests and raw prospectivity pass, so

$$E_s(X) = \{0, 19, 25, 80\} \neq E_{26}(X) = E_{70}(X) = X. \quad (145.1)$$

Directly on X , the sole near-phase key flips gap five to six, preserving the same singleton typed candidate set:

$$P_s(X) = \{0, 18, 24, 80\} \neq P_{26}(X) = P_{70}(X) = X. \quad (145.2)$$

On $E_s(X)$ the sole near-phase key at distance 19 flips gap six back to five. Hence

$$F_s(X) = \{0, 19, 24, 80\}, \quad F_{26}(X) = F_{70}(X) = X. \quad (145.3)$$

These are direct endpoint-word deductions on a four-particle malformed input, not simulator outputs or mere differences between written definitions.

145.2 The smaller phase exclusion also changes behavior

Let $Y = \{0, 5, 200, 205\}$. There are exactly two free O phase keys, at 0 and 200. No triple can use a singleton within 29 of either head; the clusters each have two particles, separated by more than $2b$. Each individual phase flip preserves both keys and creates no cross-cluster close pair. New $H_P = 4b = 148$ permits both, whereas old $H_P = 2(b + 3b + 1) = 298$ rejects both. Every edge version rejects its two raw anchors, since $200 \leq 417 \leq 834$. Therefore

$$F_s(Y) = P_s(Y) = \{0, 6, 200, 206\}, \quad F_{26}(Y) = F_{70}(Y) = Y.$$

[write] E_{26} , P_{26} , F_{26} are Report 26's $E_{||}$, $P_{||}$, F_{new} , and E_{70} , F_{70} are Report 70's E_* , F_* . Report 27's evaluator and Report 28's certificates, which are for F_{26} , therefore evaluate and certify neither F_s nor F_* .

146 Why the retained pair padding matters

The following counterexamples concern failure of admissible simulation agreement with the *unchanged literal families and stated selector*. They require an enabled nonzero branch, hence an admissible moving corridor. They are not universal lower bounds over other encodings or rules, nor failures of full-shift involutivity.

Suppose the pair exactness integer were $B_2 \leq \ell < L$. In an enabled corridor take a plus head at departure distance $t = L$. The intended behind triple remains recognized at departure marker A . The free pair at head key x would now also pass exactness: A lies at distance $L > \ell$, while the other markers are remote. These are distinct raw anchors only L apart, within H . They conflict, so selection rejects the intended travel edge. The corridor contains this position because $N \geq Z > 2(L + 1)$.

Conversely, for any integer $\ell > L$, a plus head at departure distance $L + 1$ loses its intended free pair because the departure marker is inside the enlarged window. No behind-plus template extends past L , the other markers cannot provide a triple, and this is not an interaction wall.

Thus no edge type performs the intended free step. This persists even if enlarged-read locality bookkeeping is adjusted; retaining the old read bound when $\ell > b$ would be an additional locality error, unnecessary to the counterexample.

Exactly L is therefore the retained pair padding for this particular literal free/contact subdivision and selector. Redesigned contact ranges, recognition, encoding or selection are outside the claim.

147 Evidence, provenance and reproducible presentation

This companion article is based on the complete short-exactness proof and a fresh independent mathematical/static audit. The audit’s disposition is pass without mathematical correction. It checks the endpoint families, all-input status arguments, every signed boundary and old-marker case, malformed inequalities, affine substitutions and same-rule support refinement. It limits the pair-padding examples as in Section 146. These inputs are retained in full as inert evidence, with their source bytes, modes and nanosecond modification times preserved.

The externally supplied source manifest and proof pins are respectively

```
72f8024c49a364bfd26ccc9f0e89d128e4390c1bf045477aa3d584bba3eaea31
d2229bcae33affd26eefb513e942891edf1d09c8b8d9e32ab1f8220fd2ea9962.
```

The fresh independent audit manifest and audit-text pins are respectively

```
befa9a29fd2c31a93bfecdda67794a830e5ebbedc71c4c76b85fd3fb4f161614
09cca9ebcc31338b083a502f3345e7825869fcdf043350ce16421d4ef5f7b9bd.
```

The sealed Report 70 predecessor has manifest and PDF pins

```
0f96f8aa8bba3eb581070a3db48798ecba3ef0c49daa47a124e62b6754ee30e1
d061fda50f26cad711e4f20d028d126d1ce232469425b0745728eed014895010.
```

Report 70 is a separate, unchanged companion release. Its radius-two appendix is neither duplicated nor a premise of the present theorem. The inherited proof/audit dependencies and source archive provenance remain available in the evidence, including the Report 26 archive pin

```
20a23b1ee22aed461942da4def6bade2bd55d777fce1fe49b6ff83248b7442e4.
```

[write] The two Report 70 pins are those of the delivered Report 70 archive, whose article is printed before this Report; the Report 26 archive pin is that of the committed archive of `db37d18c8`. The source packet, its seal and the fresh audit are shipped with the prefixes `71-short-windows-pp-`, `71-short-windows-seal-` and `71-short-windows-audit-fresh-`; the predecessor proof packet and the “accepted” audits consist of byte copies of Report 70’s files, of Report 26’s archive and of files of Parts II, IV and VI, which are shipped once, and of checksum ledgers, which are not shipped (README).

147.1 Historical preservation is qualified precisely

The short-exactness source packet observed its frozen inputs at 15:24:54.675013 and 15:32:53.244416 UTC on 4 October 2026. Its 49 frozen entries matched, but the separately observed live Report 70 tree was *not* static: 49 new file/directory entries appeared and the existing `qa` directory’s modification time changed. The source packet identifies these as concurrent activity, not edits made by its worker. Its mathematical inputs were the pinned frozen sources, not the changing Report 70 QA tree.

The later independent audit covered 117 frozen objects and 204 live Report 70 objects. Its observations at 15:39:29.924008 and 15:48:32.615844 UTC had no endpoint differences. That later equality does not erase the earlier live-tree qualification. The present release’s own source-origin and before/after records separately document its preservation interval. All such comparisons are endpoint observations, not atomic snapshots or proof that a tree remained continuously static between observations. Access times are outside this preservation scope and may change through reading.

147.2 What was executed and what authentication establishes

The mathematical proof is not proof-assistant formalized. No author or upstream scientific program, source interpreter, CA/physical/trajectory simulator, saved schedule or Lean was run for this article. Scientific helper programs and receipts are retained as inert source evidence. The new executable work is limited to inspected presentation, file/JSON/hash authentication, archive operations and release-tool tests. Reproducible typesetting is evidence about presentation, not a theorem prover.

The owned presentation tools adapt fully inspected Report 70 tools. The flattened `Report71.tex` matches the modular manuscript. Builds use an allowlisted environment, fresh local TeX format and font map, shell escape disabled, three TeX passes and the union of the format and every pass’s recorded inputs. Locked replay checks interpreter/executable hashes and all recorded system TeX/font inputs and selected maps. Its scope does not inventory shared libraries, Python’s standard library or the operating system. Every rendered page receives visual review in addition to automated PNG structural validation.

The release manifest records the exact descendant inventory, bytes, modes and nanosecond modification times. Authentication uses an external manifest digest and rejects unsafe paths, extra entries, symlinks, special files and hardlinks. Deterministic ZIP creation is paired with an authenticated extractor that restores descendant modes and modification times. The manifest file’s metadata and newly created top-level root’s metadata lie outside that inventory boundary; access times and creation/change times are not transported as integrity claims. Digests authenticate identity relative to supplied pins; they are not digital signatures, proof of authorship or a proof of the missing source theorem.

The README contains the final manuscript and dependency pins, authentication and locked-build commands, and deterministic archive/extraction instructions. The final release retains manuscript review, presentation-tool review and release-check receipts. No prior malformed evaluator, arithmetic certificate or universal-source execution is claimed as new evidence. No external-priority survey or radius minimality result is asserted.

148 Conclusion

Triple exactness can be shortened to its write radius in both newly defined conservative blocks, because the literal admissible geometry leaves every other marker beyond even the old window. Pair padding cannot be changed in the same way without disrupting the unchanged subdivision. The new full-shift extension preserves every admissible forward and inverse step, with main sufficient bound $76D + 105 + 3J$ and same-rule support refinement $76D + 104 + 3J$. The conditional universal numerical value is 38 722 712; its inherited-source qualifications remain essential.

149 Further questions and research for Part VII

[write] (Third write, 4 October 2026.) An unproved claim is never dropped from this collection: it is kept as an open item, with its source, the argument offered and what is missing. This section collects such claims of Reports 70 and 71, of their shipped reviews, of the programme’s intake (a9ab9a698) and of this write. Report 70’s own directions are in Section 139; Report 71 states none. No claim of the two Reports was found false.

1. *Report 68, never delivered* (a numbered remark rather than a question). Report 70 states (Section 138, and 70-two-scale-qa-pred-contract-README.md) that its presentation tools “adapt the fully inspected accepted Report 68 presentation contract”. Report 68 was never delivered to ProveIt, so its inspection and acceptance cannot be checked; that part of the sentence stays unverified. What the shipped files allow was checked for this write: reversing the shipped adaptation diff (data/70-two-scale-qa-pred-contract-ADAPTATION.diff) on Report 70’s build_report70.py, release70.py and selftest70.py gives exactly Report 68’s three tools as shipped in signal-machine-collision-certificates (code/38-low-arity-qa-pred-tools-*), with “68” renamed “70”. Nothing mathematical in this Part rests on Report 68. Missing: Report 68’s archive.
2. *Priority of the radius-two classification*. Report 70 claims none for Theorem 136.1: its search found the bibliographic record of Schranko and de Oliveira [35] but not the text, so whether that paper already classifies the conservative injective five-input binary rules is open. The finite certificate itself was rerun on copies for this write (428 conservative tables, 423 explicit collisions, five shifts; README, “Rerunning the checks”). Missing: the 2011 paper and a literature search.
3. *The conservation criterion in the literature* (this write). Lemma 136.2 is proved in Report 70 and is the radius-two analogue of SMC source 16’s elementary-rule certificate. Whether its exact de Bruijn-potential form is stated in the number-conserving literature cited in Section 98 (Boccaro–Fukš, Durand–Formenti–Róka, Moreira) was not checked; the note at the start of Section 136 attributes nothing. Missing: inspection of those papers.
4. *True and minimal radii*. Both Reports prove sufficient bounds only: $108D + 149 + 3J$ for F_* , and $76D + 105 + 3J$ (or $76D + 104 + 3J$) for F_s . The actual radii of these rules are not computed, and no lower bound is known for two-involution extensions of these admissible dynamics beyond the radius three that Report 70’s classification gives for any computationally universal binary, number-conserving, injective rule (Section 136). Report 71 shows that the pair exactness cannot be moved from L for the unchanged templates and selector (Section 146), so a further reduction needs redesigned contact ranges, recognition, encoding or selection (Report 71) or type-dependent geometric separation (Report 70, Section 139, first direction). Missing: any lower-bound argument for this encoding, and any computation of a true radius.
5. *Report 26’s own rule at isolation $b + r$* . The M2 placement’s randomized probe (dated note in Section 100) found no counterexample when Report 26’s typed-key selection isolates at $b + r$ or $2(b + r) - 1$ instead of $2(b + r)$. Reports 70 and 71 do not decide this, because each changes the recognition and hence the rule. Open: is Report 26’s construction still an involution, and still in agreement on admissible states, with isolation $b + r$? Missing: a proof or a counterexample.
6. *Evaluators and arithmetic certificates for F_* and F_s* (Report 70, Sections 137 and 139; the programme’s intake). Neither rule has an implemented evaluator or an arbitrary-input arithmetic certificate. Report 27 (Theorem 101.1) and Report 28 (Theorem 111.1) supply both for F_{26} by complete raw-key discovery and prospective rediscovery; the natural route

for F_* and F_s is the same with recognition by anchor, mutual guarded uniqueness and, for F_s , the shorter windows. Missing: the construction, its completeness proof on malformed multihead supports, and its cost audit.

7. *Independent checking of the main proofs.* The programme's intake read the statements and resource sections but did not re-audit Report 70's full-shift theorem or Report 71's recognition and geometry proof; the audits that accept them (`70-two-scale-audit-static-FRESH_AUDIT.md`, `70-two-scale-pp-INDEPENDENT_AUDIT.md`, `71-short-windows-audit-fresh-AUDIT.md`) were made by the same pipeline. No simulator of either rule was run; the malformed witnesses of Sections 135 and 145 are finite hand calculations, rechecked by hand for this write. Missing: an independent audit of Theorems 128.1, 127.1 and 140.1, or a formalization of the anchor and signed-corridor arguments (Report 70's suggestion).

A Provenance of the merge

[write] This report was built from **six manuscripts**, all delivered on 3 October 2026 as batch 82 of ProveIt's incoming reports (cluster M1) and arriving unchanged in commit `db37d18c8`. Their files were placed by commit `bca6383e9`; the text was merged by the write that added this appendix. [write] (Second write, 3 October 2026.) Four further manuscripts of the same batch (cluster M2), arriving in the same commit, were added as Parts V–VI; their files were placed by commit `7d2b1b245`. The report is now built from **ten manuscripts**. [write] (Third write, 4 October 2026.) Two manuscripts of batch 91 (cluster W4), arriving unchanged in commit `0d7f51c44`, were added as Part VII; their files were placed by commit `22a8ca89e`. The report is now built from **twelve manuscripts**.

Report	archive (batch-82 number unless stated)	file prefix	contribution and printed place
15	Reversible_Binary_Five_Particle_Package.zip (15)	15-rev-five-	base: <code>report15.tex</code> with its three <code>\input</code> files (inlined here) and figure; Sections 4–17
14	five-particle-binary-portable.zip (23)	14-five-binary-	<code>report/five-particle-binary-compile.r.tex</code> and its figure; Sections 18–32
16	Literal_Universal_Reversible_Source_Package.zip (13)	16-literal-source-	<code>report16.tex</code> with its ten section files (inlined); Part II
17	Reversible_Startup_Optimization_Package.zip (16)	17-startup-	<code>report17.tex</code> ; first half of Part III
18	Cellular_Clock_Domination_Package.zip (04)	18-cell-clock-	<code>report18.tex</code> ; second half of Part III
19	Exact_Lazy_Reversible_CA_Evaluator_Package.zip (09)	19-lazy-eval-	<code>report19.tex</code> ; Part IV
20	Event_Budgeted_Quartic_Certificates_Package.zip (08)	20-event-budget-	<code>report20.tex</code> ; Part V, Sections 77–88
26	Two_Parallel_Conservative_Involutions_Package.zip (20)	26-parallel-involutions-	<code>report26.tex</code> ; Part VI, first: Sections 89–100
27	Sparse_Parallel_Particle_Evaluation_Package.zip (18)	27-sparse-parallel-	<code>report27.tex</code> ; Part VI, second: Sections 101–110
28	Canonical_Parallel_Quartic_Certificates_Package.zip (03)	28-parallel-quartic-	<code>report28.tex</code> ; Part VI, third: Sections 111–126
70	Two_scale_recognition_radians_sources.zip (batch 91: 22)	70-two-scale-	<code>Report70.tex</code> , the flattened form of its seven <code>manuscript/</code> modules; Part VII, first: Sections 127–139
71	Shorter_exactness_window_s_sources.zip (batch 91: 23)	71-short-windows-	<code>Report71.tex</code> , the flattened form of its six modules; Part VII, second: Sections 140–148

Pins. No manuscript names a ProveIt commit, so none has a repository pin; each is dated 3 October 2026. The Reports pin one another by SHA-256: Report 16 pins Report 15’s compiler (`f95a6028c9d3...`); Report 17 pins Report 16’s source (`38c586706fa1...`); Reports 18 and 19 pin Report 17’s source (`fa61d06178d7...`); Reports 14–16 pin Report 12 (`803bf0c4194b...`, the L^AT_EX source of SMC source 17). The file prefixes are the series’ Report numbers rather than batch numbers, because batch numbers 13, 15 and 16 coincide with other Reports’ numbers. Reports 20 and 26–28 name no ProveIt commit either. They pin Report 15’s compiler (`f95a6028c9d3...`), Report 19’s evaluator (`42e8aa65c05f...`, Reports 20, 27 and 28) and Report 17’s source (`fa61d06178d7...`, Reports 20 and 27); Report 27 pins Report 26’s scientific manifest (`141f1ecf4e34...`), and Report 28 pins Report 26’s evaluator (`2c8b6396...`) and Report 27’s (`785082d12162...`). Reports 70 and 71 name no ProveIt commit either; both are dated 4 October 2026 and name no author. Report 70 pins Report 26’s archive (`20a23b1ee22a...`, the committed archive of `db37d18c8`), Report 19’s universal receipt and the absent source table `fa61d06178d7...`; Report 71 pins Report 70’s release manifest (`0f96f8aa8bba...`) and PDF (`d061fda50f26...`), both equal to the delivered Report 70 archive’s files, and Report 26’s archive.

Where the merge had to choose.

1. *Order and grouping.* Report 15 (the base, and the compiler used by Reports 16–19) is printed before Report 14; Reports 17 and 18 share Part III. Each Report keeps its own sections, numbered continuously; its abstract and any front paragraphs are printed at its opening.

2. *Duplicated results.* Report 15’s threshold implies Report 14’s (a globally reversible number-conserving binary automaton is a binary conservative automaton); both are printed, since the compilers differ, and Report 14’s corollary carries a pointer. Re-proofs of results printed in SMC and QOC are kept as second presentations with pointers (Section 1.6); nothing was deleted. Report 19’s restatement of Report 15’s ledger and the four occurrences of the compiler clock carry pointers.
3. *Bibliography.* The six bibliographies are merged into one. Entries for the companion Reports 12, 15, 16 and 17 point to their printed places (SMC Part VI, or a Part of this report); entries for bundled notes point to the shipped files. Bennett [8] is added for the clean-target wrapper.
4. *Inputs.* Report 15’s `certificate-section.tex`, `clean-target-section.tex` and `verification-details.tex`, and Report 16’s `sections/*.tex`, are inlined; Report 19’s appendix “Byte identities” becomes an ordinary section (Section 76), and its prose reference “Section 2” became a cross-reference.
5. *Macros.* The union of the six preambles is used; Report 14’s `\enc` (printing “Enc”) is renamed `\Enc` internally so that Reports 15–19 keep `\enc` (printing “enc”). No printed symbol changed.
6. *Parts V–VI: order and grouping* (second write). Report 20, the only Report on the old ordered rule, forms Part V; Reports 26–28, on the new parallel rule, share Part VI in Report order. The M2 dossier proposed four Parts; one Part per rule was chosen instead. Each Report keeps its own sections; its abstract and a dated note open it, and its table of contents is dropped.
7. *Parts V–VI: duplicates.* Report 27’s Lemma 103.1 restates Report 26’s Lemma 90.1 with a condensed proof; it is kept, marked as a restatement, together with Report 27’s rederived radii, and Report 26’s malformed witness, which Reports 27 and 28 repeat, carries pointers. The arithmetic layer of Reports 20 and 28 and Report 28’s sorter are kept as second presentations credited to CDC and SMC Part IV. Reports 20 and 28 share a compiler pattern but prove different theorems for different rules; both are printed, with a note on what they share. Nothing was deleted.
8. *Parts V–VI: bibliography.* Reports 26 and 27’s entries for Boyle–Lind–Rudolph, Toffoli–Margolus, Gajardo–Kari–Moreira and Maldonado–Moreira–Gajardo are merged (Report 26’s URLs and Report 27’s DOIs together); Report 20’s “Report 8” entry points to SMC Part IV and its “Report 19” entry to Part IV; entries for Reports 26 and 27, which Report 28 names without citing, are added; seven entries on partitioned and number-conserving automata are added for the lineage note in Section 98.
9. *Parts V–VI: macros and symbols.* Report 26’s `\N` ($\mathbb{Z}_{\geq 0}$) and Report 28’s `\N` ($\mathbb{N}_0$) are `\Nge` and `\Nzero` internally, Report 27’s `\code` and `\C` are `\file` and `\calC`; these print as delivered. One printed symbol changed: Report 20’s height bound H_* is printed Ht_* , because Report 28 uses H_* for an isolation distance.
10. *Parts V–VI: appendices and receipts.* The appendices of Reports 27 and 28 are ordinary sections. A footnote in Report 28’s Section 123 qualifies its statement that all receipts retain pinned identities (two archived supplement receipts predate the final compiler).
11. *Part VII: placement and order* (third write). Reports 70 and 71 answer Report 26’s first question and keep its templates and admissible graph, so they form a Part of this report rather than a new one; Report 70 is printed first because Report 71 pins it. The file prefixes are the Report numbers, as in Parts I–VI.

12. *Part VII: duplicates.* Report 71 restates Report 70’s source interface, templates and admissible graph with its new exactness predicate; both are printed. Report 70’s phase lemma restates Report 26’s phase block, and its source, template, clock and clean-target sections restate Reports 15 and 26; each carries a pointer. Nothing was deleted. Both packages carry byte copies of Report 26’s archive and files, of Report 19’s universal receipts and of Report 16’s copy of Report 15’s compiler proof, which are shipped once (README).
13. *Part VII: unproved claims.* Claims of the two Reports, of their shipped reviews, of the programme’s intake and of this write that this report does not prove are not dropped; they are collected as open items, with source, argument and what is missing, in Section 149. None of them was found false; the one claim resting on the undelivered Report 68 concerns presentation tooling only and was checked as far as the shipped files allow.
14. *Part VII: labels, macros and bibliography.* The eleven label names the two Reports share are kept apart by the sub-prefixes `fpa:tr:` and `fpa:sw:.` Report 70’s `\Fold` (printing F_{26}) clashes with Part VI’s `\Fold` (F_{old}) and is `\Ftwosix` here, with `\Eold`, `\Pold` renamed `\Etwosix`, `\Ptwosix` to match; `\N` is `\Nge`; no printed symbol changed. Report 70’s entries for Boyle–Lind–Rudolph, Toffoli–Margolus, Maldonado–Moreira–Gajardo, Gajardo–Kari–Moreira and Morita are the existing entries; Fuk’s–Sullivan and Schranko–de Oliveira are added. Report 71 has no bibliography. Both tables of contents are dropped, and the Report openings print the delivered date, 4 October 2026.

Correction of a citation (dated 3 October 2026). Reports 14, 16, 17 and 18 cite Neary and Woods as *Fundamenta Informaticae* 91 (2009), 105–126, with DOI 10.3233/FI-2009-0008, and Reports 14 and 16 call the page range 123–144 incorrect. The DOI 10.3233/FI-2009-0008 belongs to a different paper (Han, Salomaa and Wood, *Nondeterministic state complexity of basic operations for prefix-free regular languages*, *Fundamenta Informaticae* 90 (2009), 93–106). The publisher’s record of the Neary–Woods paper is volume 91, issue 1, pages 123–144, DOI 10.3233/FI-2009-0036 (checked against Crossref on 3 October 2026), as quadratic-orthant-certificates already cites it; 105–126 is the pagination of the author-hosted PDF, which the author-deposited archive flags as incorrect. The merged bibliography entry [1] follows the publisher; the page locators “printed p. 121” and “printed p. 123” in the Reports refer to the author-hosted PDF. The Reports’ sentences are kept, each followed by a dated note; the delivered audit and provenance files that repeat the old DOI are byte-identical and listed in the README.

Labels. The six manuscripts carry 194 labels (Report 14: 59, 15: 40, 16: 23, 17: 30, 18: 31, 19: 11); all are kept with their prefixes, and the write added the labels named in Section 2. The four manuscripts of Parts V–VI carry 82 labels (Report 20: 22, 26: 33, 27: 6, 28: 21); all are kept with their sub-prefixes, and the second write added 13 (Section 2). The two manuscripts of Part VII carry 84 labels (Report 70: 49, Report 71: 35), eleven of them with the same delivered names; all are kept with their sub-prefixes, and the third write added 6 (Section 2).

References

- [1] T. Neary and D. Woods. *Four Small Universal Turing Machines*. Fundamenta Informaticae **91**(1) (2009), 123–144. DOI: [10.3233/FI-2009-0036](https://doi.org/10.3233/FI-2009-0036). Primary author-hosted copy: <https://dna.hamilton.ie/assets/dw/NearyWoods-FI09.pdf>, paginated 105–126; the author-deposited archive notes that pagination as incorrect, and the publication details here follow the publisher. [write] Corrected 3 October 2026: Reports 14, 16, 17 and 18 printed “91 (2009), 105–126, doi:10.3233/FI-2009-0008”; that DOI is a different paper (Appendix A). Locators used by the Reports, in the author PDF’s pagination: finite input, Eq. (3), Definition 3.1 and Table 1 (Reports 14, 16); literal machine, Section 3.5 and Table 16, printed p. 121 (PDF p. 17); halt display and prose discrepancy, printed p. 123 (PDF p. 19). Report 18: “finite-input theorem and Table 16 are inherited dependencies.”
- [2] K. Morita. *Universality of a reversible two-counter machine*. Theoretical Computer Science **168**(2) (1996), 303–320. doi:10.1016/S0304-3975(96)00081-3. Used for Definitions 2.1–2.3 and Theorems 4.1–4.2 (Report 15), Theorems 3.1 and 4.1 (Report 17: “the exact bundled graphs are separately proved”), and as method attribution (Report 18: “the actual graphs are independently reconstructed”).
- [3] K. Morita and K. Imai. *Number-conserving reversible cellular automata and their computation-universality*. RAIRO Theoretical Informatics and Applications **35**(3) (2001), 239–258. [Primary publisher record](#). In particular Lemma 3.5 and Theorem 3.6.
- [4] K. Morita. *Universality of One-Dimensional Reversible and Number-Conserving Cellular Automata*. 2012. [arXiv:1208.2760](https://arxiv.org/abs/1208.2760). [write] Report 70 cites it as EPTCS 90 (2012), 142–150, for Morita’s 96-state four-neighbour reversible number-conserving automaton, a model distinct from the binary five-particle compiler.
- [5] A. Alhazov and K. Imai. *Particle Complexity of Universal Finite Number-Conserving Cellular Automata*. CANDAR 2016, 209–214. doi:10.1109/CANDAR.2016.0045; [publisher abstract](#). The five-particle numerical upper bound is attributed at abstract level; the full construction was not available to, or inspected by, the reviews of Reports 14 and 15.
- [6] G.-T. Kong. *A hierarchical structural interpretation of 1-dimensional 2-state number conserving cellular automata*. Doctoral thesis, Hiroshima University, June 2021, Section 2.4. [Primary thesis](#). The lower-particle statements cited in Report 12 concern binary NCCAs.
- [7] K. Imai. *Structure and computing efficiency of number-conserving cellular automata*. KAKEN project 17K00015 final research report, 25 June 2021, PDF pp. 3–4. [Official research report](#). Its short decidability statement does not specify every decision problem or all the hypotheses of Report 12’s theorem.
- [8] C. H. Bennett. *Logical Reversibility of Computation*. IBM Journal of Research and Development **17**(6) (1973), 525–532. <https://doi.org/10.1147/rd.176.0525>. [write] Added in the merge, for the clean-target wrapper of Section 13.
- [9] S. Basu. *Algorithms in Real Algebraic Geometry: A Survey*. Author-hosted survey, Theorem 2.1, PDF page 5. https://www.math.purdue.edu/~sbasu/raag_survey2011_final-sep4-2014.pdf.
- [10] *Four mass units and exact reachability*. Report 12 of the series, 3 October 2026: the four-mass decision theorem and binary timing certificate, an AI-assisted research report, not a refereed publication or proof-assistant formalization; Report 14: “Independently inspected mathematical decision theorem and binary timing certificate . . . No primary publication or priority status is asserted.” [write] It is printed as source 17 (Part VI) of the collection

- report `signal-machine-collision-certificates` (its Theorem `smc:fm:thm:main` and Corollary `smc:fm:cor:numerical`); the byte-identical companion copies of its L^AT_EX source, PDF and figure that Reports 14, 15 and 16 delivered are not shipped here. Reports 14–16 use it as the four-mass lower bound, “with its stated model and qualifications” (Report 16).
- [11] *Globally reversible binary computation with five particles*. Report 15, 3 October 2026; printed as the first half of Part I. Its compiler is shipped once, as `code/15-rev-five-compiler-reversible_binary.py` (SHA-256 `f95a6028c9d3...`); Reports 16–19 delivered byte-identical copies of it, of its source schema and compiler proof, and Reports 16 and 19 copies of its PDF, which are not shipped. Report 16: “The universal source there was existential; the present report materializes it separately.”
 - [12] *A literal universal reversible two counter source*. Report 16, 3 October 2026; printed as Part II. Report 17: “Frozen predecessor construction; the exact source hash is recorded in Section 42. This report is a separate follow-up, not a replacement edition.”
 - [13] *Startup history in a literal reversible source*. Report 17, 3 October 2026; printed as the first half of Part III. Report 19 uses the exact source identified in Section 76.
 - [14] *Direct primitive shared-offset finite-horizon quartic certificates*. 3 October 2026. Report 18’s “inherited proof, especially Sections 5–7”; it is Report 16’s shared-offset certificate (Section 39.4), shipped as `16-literal-source-primitive-certificates-PROOF.md` (Report 18 delivered a byte-identical copy as `certificate-reference/SHARED_OFFSET_PROOF.md`).
 - [15] *Cellular-clock domination for the fixed history-orientation templates*. 3 October 2026. Report 18’s bundled producer proof and independent audit, shipped as `18-cell-clock-cellular-clock-CA_CLOCK_DOMINATION.md` and `18-cell-clock-cellular-clock-audit-AUDIT.md`.
 - [16] *Lazy exact evaluation of the frozen ordered binary CA*. Report 19’s implementation proof, shipped as `19-lazy-eval-PROOF.md`.
 - [17] *Exact lazy finite support evaluation of the frozen reversible binary CA*. Report 19, 3 October 2026; printed as Part IV (Theorem 65.1). Report 20: “The bundled dependency proof is `reproducibility/dependencies/report19-proof.md`; evaluator and frozen compiler bytes are locally pinned.” [write] That bundled proof is a byte copy of `19-lazy-eval-PROOF.md` and is not shipped again.
 - [18] Report 8. *Canonical natural comparison and sum-of-squares arithmetic, with signed constant division in the semilinearity treatment*. Report 20 bundled the reference sources `report8-core.tex` and `report8-semilinearity.tex` under `reproducibility/dependencies/`, “used as prior arithmetic constructions, with the required gadgets proved again here”. [write] Report 8 is source 13 (Part IV) of `signal-machine-collision-certificates`; the two excerpts are its own text (the opening sections of Part IV, and its Sections `smc:sl:sec:semilinearity` and `smc:sl:sec:presburger-quartic`, differing only in the macro for the naturals) and are not shipped here. The same gadgets are in `canonical-diophantine-certificates` (note after Lemma 79.1).
 - [19] *Companion finite-schema proof and executable evidence*, bundled with Report 20: `reproducibility/PROOF.md`, sample coefficient and witness files, scheduler tests and independent audit. [write] Shipped as `20-event-budget-PROOF.md`, `data/20-event-budget-example-polynomial.json`, `data/20-event-budget-example-witness.json`, `code/20-event-budget-scheduler-reference.py` and `20-event-budget-audit-INDEPENDENT_AUDIT.md`.

- [20] *Two parallel conservative involutions*, subtitle *A full shift replacement for the five particle compiler*. Research report 26, 3 October 2026; printed as the first Report of Part VI. Report 27: frozen 3 October 2026; its bundled scientific proof is titled *Two parallel involutions for the binary five-particle construction*; “Proof and independent lemma/preservation audits are reproduced unchanged under `scientific/PARALLEL_*.md`; provenance in `references/report26-provenance.json`.” [write] Those copies are byte-identical to Report 26’s files and are shipped once, as `26-parallel-involutions-PROOF.md`, `26-parallel-involutions-audit-lemma.md`, `26-parallel-involutions-audit-preservation.md` and `data/26-parallel-involutions-delivery-provenance.json`. [write] The entry is added for Report 28, which names Reports 26 and 27 without bibliography entries.
- [21] *Sparse evaluation of a two involution particle cellular automaton*. Research Report 27, 3 October 2026; printed as the second Report of Part VI. [write] Added for Report 28.
- [22] M. Boyle, D. Lind and D. Rudolph. *The automorphism group of a shift of finite type*. Transactions of the American Mathematical Society **306** (1988), 71–114. <https://doi.org/10.1090/S0002-9947-1988-0927684-2>; author copy <https://sites.math.washington.edu/~lind/Papers/AutomorphismGroupSFT.pdf>. Report 26: Section 2 and Theorem 2.6.
- [23] T. Toffoli and N. H. Margolus. *Invertible cellular automata: a review*. Physica D **45** (1990), 229–253, Sections 5.3 and 5.5 (Report 27: especially Section 5.3). [https://doi.org/10.1016/0167-2789\(90\)90185-R](https://doi.org/10.1016/0167-2789(90)90185-R); author PDF <https://people.csail.mit.edu/nhm/ica.pdf>.
- [24] A. Gajardo, J. Kari and A. Moreira. *On time-symmetry in cellular automata*. Journal of Computer and System Sciences **78** (2012), 1115–1126, Proposition 1. <https://doi.org/10.1016/j.jcss.2012.01.006>; author preprint <https://www.ci2ma.udec.cl/pdf/pre-publicaciones2/2011/pp11-28.pdf>.
- [25] D. Maldonado, A. Moreira and A. Gajardo. *Universal Time-Symmetric Number-Conserving Cellular Automaton*. AUTOMATA 2015, Lecture Notes in Computer Science **9099**, 155–168, especially Proposition 2. https://doi.org/10.1007/978-3-662-47221-7_12; author preprint <https://www.ci2ma.udec.cl/pdf/pre-publicaciones/2015/pp15-18.pdf>.
- [26] D. E. Knuth. *The Art of Computer Programming*, volume 3, *Sorting and Searching*, second edition, Addison-Wesley, 1998, Section 5.3.4, “Networks for Sorting”. ISBN 0-201-89685-0. Author’s page: <https://www-cs-faculty.stanford.edu/~knuth/taocp.html>.
- [27] K. Morita and M. Harao. *Computation universality of one-dimensional reversible (injective) cellular automata*. Transactions of the IEICE **E72**(6) (1989), 758–762. [write] Added in the write (lineage note in Section 98); not inspected.
- [28] J. Kari. *Representation of reversible cellular automata with block permutations*. Mathematical Systems Theory **29**(1) (1996), 47–61. DOI [10.1007/BF01201813](https://doi.org/10.1007/BF01201813) (Crossref, checked 3 October 2026). [write] Added in the write; Report 26’s audit checked it at abstract level only.
- [29] J. Kari. *On the circuit depth of structurally reversible cellular automata*. Fundamenta Informaticae **38** (1999), 93–107. DOI [10.3233/FI-1999-381208](https://doi.org/10.3233/FI-1999-381208) (Crossref, checked 3 October 2026). [write] Added in the write; Report 26’s audit checked it at abstract level only.
- [30] J. Durand-Lose. *Representing reversible cellular automata with reversible block cellular automata*. Discrete Models: Combinatorics, Computation, and Geometry (DM-CCG 2001), Discrete Mathematics and Theoretical Computer Science Proceedings **AA** (2001), 145–154. [write] Added in the write; not inspected.

- [31] N. Boccara and H. Fuk s. *Cellular automaton rules conserving the number of active sites*. Journal of Physics A: Mathematical and General **31**(28) (1998), 6007–6018. DOI [10.1088/0305-4470/31/28/014](https://doi.org/10.1088/0305-4470/31/28/014) (Crossref, checked 3 October 2026). [write] Added in the write; not inspected.
- [32] B. Durand, E. Formenti and Zs. R ka. *Number-conserving cellular automata I: decidability*. Theoretical Computer Science **299** (2003), 523–535. DOI [10.1016/S0304-3975\(02\)00534-0](https://doi.org/10.1016/S0304-3975(02)00534-0) (Crossref, checked 3 October 2026). [write] Added in the write; not inspected.
- [33] A. Moreira. *Universality and decidability of number-conserving cellular automata*. Theoretical Computer Science **292**(3) (2003), 711–721. DOI [10.1016/S0304-3975\(02\)00065-8](https://doi.org/10.1016/S0304-3975(02)00065-8) (Crossref, checked 3 October 2026). [write] Added in the write; not inspected.
- [34] H. Fuk s and K. Sullivan. *Enumeration of number-conserving cellular automata rules with two inputs*. 2007 preprint, Theorem 1 and Section 2. <https://arxiv.org/pdf/0711.1349>. Report 70: the count 428 of conservative five-input binary tables; its primary-source check (70-two-scale-qa-PRIMARY_SOURCE_CHECK.md) records the counts 1, 2, 5, 22, 428 on PDF page 2. [write] Added from Report 70’s bibliography.
- [35] A. Schranko and P. P. B. de Oliveira. *Derivation and representation of one-dimensional, reversible, number-conserving cellular automata rules*. Journal of Cellular Automata **6**(1) (2011), 77–89. Report 70: bibliographic record checked on the author’s publication list, <https://professor.mackenzie.br/pedrob/RESEARCH/PAPERS/publications-Themes.html>; full text not obtained, so no theorem or priority is attributed to it. [write] Added from Report 70’s bibliography.
- [36] *Independent audit of the lazy ordered reversible evaluator*. Report 19’s audit narrative, addendum, executable checks and hash-linked receipts, shipped as 19-lazy-eval-audit-design.md, 19-lazy-eval-audit-addendum.md and the 19-lazy-eval- files of code/ and data/.